

# MLSys-im Tutorial — Module 4

## Design Space Exploration & Synthesis

---

Vijay Janapa Reddi

Conference Tutorial

Harvard University

# Roadmap: Conference Tutorial

---

| Module          | Topic                                           | Status                |
|-----------------|-------------------------------------------------|-----------------------|
| Module 1        | Foundations & Architecture                      | ✓ Done                |
| Module 2        | Advanced Single-Node Analysis                   | ✓ Done                |
| Module 3        | Scale, Dollars, and Carbon                      | ✓ Done                |
| <b>Module 4</b> | <b>Design Space Exploration &amp; Synthesis</b> | <b>← You are here</b> |

TUTORIAL | MODULE 4

# Rapid Parametric Sweeps

---

1

# The Combinatorial Explosion

---

Finding the optimal serving configuration requires testing:

$$|\text{hardware}| \times |\text{batch sizes}| \times |\text{precisions}| \times |\text{parallelism configs}|$$

This space easily exceeds  $10^4$  configurations. Because `mlsysim` uses analytical math (not cycle-accurate simulation), each evaluation takes  $< 1$  ms.

# Live Demo: Programmatic Sweeps

```
from mlsysim.engine.engine import Engine
from mlsysim.hardware.registry import Hardware
from mlsysim.models.registry import Models
import pandas as pd

model = Models.Language.Llama3_8B
hardware = Hardware.Cloud.H100

results = []
for batch_size in [1, 8, 32, 128, 256]:
    perf = Engine.solve(model, hardware, batch_size=batch_size, precision="fp16")
    results.append({
        "Batch Size": batch_size,
        "Throughput (tok/s)": perf.throughput.magnitude,
        "Bottleneck": perf.bottleneck
    })

print(pd.DataFrame(results))
```

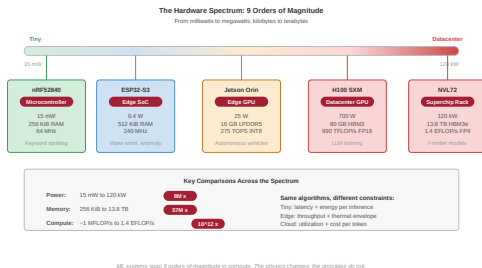
TUTORIAL | MODULE 4

# TinyML to Frontier

---

2

# Same Roofline, 9 Orders of Magnitude



| Device   | FLOPS | TDP    |
|----------|-------|--------|
| nRF52840 | 64 M  | 15 mW  |
| ESP32-S3 | 500 M | 400 mW |
| H100 SXM | 989 T | 700 W  |

Compute Range  $\sim 10^7 \times$   
Power Range  $\sim 10^{4.7} \times$

The Roofline Model is universal. The physics apply identically to a \$2 Microcontroller and a \$3M GPU Rack.

TUTORIAL | MODULE 4

# Sensitivity Analysis

---

3



# Sensitivity Analysis

```
from mlsysim.solvers import SensitivitySolver
solver = SensitivitySolver()

result = solver.solve(
    model=Models.Language.Llama3_8B, hardware=Hardware.Cloud.H100,
    precision="fp16", efficiency=0.5)

print(f"Binding Constraint: {result.binding_constraint}")
for param, sensitivity in result.sensitivities.items():
    tag = "<<<" if param == result.binding_constraint else ""
    print(f"  {param:>20}: {sensitivity:+.4f}  {tag}")
```

**The Golden Rule:** Invest in the parameter with the *largest* partial derivative. Improving a non-binding parameter yields **zero** measurable gain.

TUTORIAL | MODULE 4

# SLA-Driven Synthesis

---

4

# Live Demo: Inverting the Roofline

Instead of asking "How fast is this GPU?", what if we ask "What hardware do I need to buy to meet my 30ms latency SLA?"

```
from mlsysim.solvers import SynthesisSolver
from mlsysim.models.registry import Models
from mlsysim.core.units import Q_

solver = SynthesisSolver()
requirements = solver.solve(
    model=Models.Language.Llama3_8B,
    target_latency=Q_("30 ms"),
    batch_size=1,
    precision="fp16"
)

print(f"Required HBM Bandwidth: {requirements.required_bw.to('GB/s'):.1f}")
print(f"Required Compute: {requirements.required_flops.to('TFLOPs/s'):.1f}")
```

TUTORIAL | MODULE 4

# Capstone & Wrap-Up

---

5

# Design Challenge: The Capstone

## The Problem

**\$5M budget.** Serve Llama-3 70B at **1,000 QPS** with **<100 ms TTFT** in **two regions** (US-East + EU-West). Design the fleet.

**You must specify using `mlsysim`:**

- 1. Hardware choice:** Which GPU? How many?
- 2. Parallelism strategy:** TP  $\times$  PP?
- 3. Precision:** FP16? FP8? INT4?
- 4. Geographic placement:** Carbon impact?

# Resources & Next Steps

---

## Get Started

- `pip install mlsysim`
- GitHub: `harvard-edge/mlsysim`
- Full docs: `mlsysim.readthedocs.io`
- Code cookbook: five interactive scenarios

## The Textbook

- *Machine Learning Systems*
- Volume I: Foundations (single node)
- Volume II: Systems at Scale (fleet)
- `mlsysbook.ai`

# Key Papers

---

- Williams et al. (2009)  
Roofline Model
- Chowdhery et al. (2022)  
PaLM / MFU
- Hoffmann et al. (2022)  
Chinchilla Scaling
- Patterson et al. (2021)  
Carbon & Training
- OpenAI (2024)  
o1 / Reasoning Scaling

Use these papers to validate the assumptions behind each solver family.

# Thank you! Questions?