



Don't just import torch. Build it.

# Build Your Own ML Framework

**Prof. Vijay Janapa Reddi**  
Harvard University

*From Tensors to Systems*

A Build-It-Yourself Companion to the

**Machine Learning Systems** textbook

[mlsysbook.ai](https://mlsysbook.ai)



# Table of contents

---

|          |                                                  |           |
|----------|--------------------------------------------------|-----------|
| <b>1</b> | <b>About This Guide</b>                          | <b>1</b>  |
| <b>2</b> | <b>Welcome</b>                                   | <b>3</b>  |
| 2.1      | The Problem . . . . .                            | 3         |
| 2.2      | The Solution: AI Bricks . . . . .                | 3         |
| 2.3      | Who This Is For . . . . .                        | 4         |
| 2.4      | What You Need to Start . . . . .                 | 4         |
| 2.5      | What You Will Build . . . . .                    | 4         |
| 2.6      | How to Learn . . . . .                           | 4         |
| 2.7      | The Bigger Picture . . . . .                     | 5         |
| 2.8      | For Instructors . . . . .                        | 5         |
| 2.9      | Make It Better . . . . .                         | 5         |
| 2.10     | What's Next? . . . . .                           | 5         |
| <b>3</b> | <b>Big Picture</b>                               | <b>7</b>  |
| 3.1      | The Journey: Foundation to Production . . . . .  | 7         |
| 3.2      | Milestones You'll Unlock . . . . .               | 8         |
| 3.3      | What You'll Have at the End . . . . .            | 8         |
| 3.4      | Choose Your Learning Path . . . . .              | 9         |
| 3.5      | Expect to Struggle (That's the Design) . . . . . | 9         |
| 3.6      | Start Building . . . . .                         | 10        |
| <b>4</b> | <b>Getting Started with TinyTorch</b>            | <b>11</b> |
| 4.1      | The Journey . . . . .                            | 11        |
| 4.2      | Step 1: Install & Setup (2 Minutes) . . . . .    | 11        |
| 4.3      | macOS / Linux . . . . .                          | 11        |
| 4.4      | Windows . . . . .                                | 12        |
| 4.5      | Step 2: Your First Module (15 Minutes) . . . . . | 12        |
| 4.6      | Step 3: Your First Milestone . . . . .           | 13        |
| 4.7      | The Pattern Continues . . . . .                  | 14        |
| 4.8      | Quick Reference . . . . .                        | 14        |
| 4.9      | Module Progression . . . . .                     | 15        |
| 4.10     | Join the Community (Optional) . . . . .          | 15        |
| 4.11     | For Instructors & TAs . . . . .                  | 15        |
| <b>I</b> | <b>Foundation Tier</b>                           | <b>17</b> |
| <b>5</b> | <b>Module 01: Tensor</b>                         | <b>19</b> |
| 5.1      | Overview . . . . .                               | 19        |
| 5.2      | Learning Objectives . . . . .                    | 19        |
| 5.3      | What You'll Build . . . . .                      | 19        |
| 5.4      | API Reference . . . . .                          | 21        |
| 5.5      | Core Concepts . . . . .                          | 22        |
| 5.6      | Architecture . . . . .                           | 28        |
| 5.7      | Common Errors & Debugging . . . . .              | 29        |

|          |                                    |           |
|----------|------------------------------------|-----------|
| 5.8      | Production Context . . . . .       | 29        |
| 5.9      | Your TinyTorch . . . . .           | 30        |
| 5.10     | PyTorch . . . . .                  | 30        |
| 5.11     | Check Your Understanding . . . . . | 31        |
| 5.12     | Key Takeaways . . . . .            | 32        |
| 5.13     | Further Reading . . . . .          | 33        |
| 5.14     | What's Next? . . . . .             | 33        |
| 5.15     | Get Started . . . . .              | 34        |
| <b>6</b> | <b>Module 02: Activations</b>      | <b>35</b> |
| 6.1      | Overview . . . . .                 | 35        |
| 6.2      | Learning Objectives . . . . .      | 35        |
| 6.3      | What You'll Build . . . . .        | 36        |
| 6.4      | API Reference . . . . .            | 37        |
| 6.5      | Core Concepts . . . . .            | 38        |
| 6.6      | Production Context . . . . .       | 43        |
| 6.7      | Your TinyTorch . . . . .           | 44        |
| 6.8      | PyTorch . . . . .                  | 44        |
| 6.9      | Check Your Understanding . . . . . | 45        |
| 6.10     | Key Takeaways . . . . .            | 47        |
| 6.11     | Further Reading . . . . .          | 48        |
| 6.12     | What's Next . . . . .              | 48        |
| 6.13     | Get Started . . . . .              | 49        |
| <b>7</b> | <b>Module 03: Layers</b>           | <b>51</b> |
| 7.1      | Overview . . . . .                 | 51        |
| 7.2      | Learning Objectives . . . . .      | 51        |
| 7.3      | What You'll Build . . . . .        | 52        |
| 7.4      | API Reference . . . . .            | 53        |
| 7.5      | Core Concepts . . . . .            | 54        |
| 7.6      | Common Errors . . . . .            | 60        |
| 7.7      | Production Context . . . . .       | 61        |
| 7.8      | Your TinyTorch . . . . .           | 61        |
| 7.9      | PyTorch . . . . .                  | 62        |
| 7.10     | Check Your Understanding . . . . . | 63        |
| 7.11     | Key Takeaways . . . . .            | 65        |
| 7.12     | Further Reading . . . . .          | 65        |
| 7.13     | What's Next . . . . .              | 66        |
| 7.14     | Get Started . . . . .              | 67        |
| <b>8</b> | <b>Module 04: Losses</b>           | <b>69</b> |
| 8.1      | Overview . . . . .                 | 69        |
| 8.2      | Learning Objectives . . . . .      | 69        |
| 8.3      | What You'll Build . . . . .        | 70        |
| 8.4      | API Reference . . . . .            | 71        |
| 8.5      | Core Concepts . . . . .            | 72        |
| 8.6      | Common Errors . . . . .            | 75        |
| 8.7      | Production Context . . . . .       | 76        |
| 8.8      | Your TinyTorch . . . . .           | 77        |
| 8.9      | PyTorch . . . . .                  | 77        |
| 8.10     | Check Your Understanding . . . . . | 78        |
| 8.11     | Key Takeaways . . . . .            | 82        |
| 8.12     | Further Reading . . . . .          | 82        |
| 8.13     | What's Next . . . . .              | 83        |

|           |                              |            |
|-----------|------------------------------|------------|
| 8.14      | Get Started                  | 83         |
| <b>9</b>  | <b>Module 05: DataLoader</b> | <b>85</b>  |
| 9.1       | Overview                     | 85         |
| 9.2       | Learning Objectives          | 85         |
| 9.3       | What You'll Build            | 85         |
| 9.4       | API Reference                | 87         |
| 9.5       | Core Concepts                | 88         |
| 9.6       | Common Errors                | 92         |
| 9.7       | Production Context           | 93         |
| 9.8       | Your TinyTorch               | 94         |
| 9.9       | PyTorch                      | 94         |
| 9.10      | Check Your Understanding     | 96         |
| 9.11      | Key Takeaways                | 98         |
| 9.12      | Further Reading              | 98         |
| 9.13      | What's Next                  | 99         |
| 9.14      | Get Started                  | 99         |
| <b>10</b> | <b>Module 06: Autograd</b>   | <b>101</b> |
| 10.1      | Overview                     | 101        |
| 10.2      | Learning Objectives          | 101        |
| 10.3      | What You'll Build            | 102        |
| 10.4      | API Reference                | 103        |
| 10.5      | Core Concepts                | 104        |
| 10.6      | Production Context           | 109        |
| 10.7      | Your TinyTorch               | 110        |
| 10.8      | PyTorch                      | 110        |
| 10.9      | Check Your Understanding     | 111        |
| 10.10     | Key Takeaways                | 113        |
| 10.11     | Further Reading              | 114        |
| 10.12     | What's Next                  | 115        |
| 10.13     | Get Started                  | 115        |
| <b>11</b> | <b>Module 07: Optimizers</b> | <b>117</b> |
| 11.1      | Overview                     | 117        |
| 11.2      | Learning Objectives          | 117        |
| 11.3      | What You'll Build            | 118        |
| 11.4      | API Reference                | 119        |
| 11.5      | Core Concepts                | 120        |
| 11.6      | Production Context           | 124        |
| 11.7      | Your TinyTorch               | 125        |
| 11.8      | PyTorch                      | 125        |
| 11.9      | Check Your Understanding     | 126        |
| 11.10     | Key Takeaways                | 128        |
| 11.11     | Further Reading              | 128        |
| 11.12     | What's Next                  | 129        |
| 11.13     | Get Started                  | 129        |
| <b>12</b> | <b>Module 08: Training</b>   | <b>131</b> |
| 12.1      | Overview                     | 131        |
| 12.2      | Learning Objectives          | 131        |
| 12.3      | What You'll Build            | 131        |
| 12.4      | API Reference                | 133        |
| 12.5      | Core Concepts                | 134        |
| 12.6      | Production Context           | 141        |

|                                                             |                |
|-------------------------------------------------------------|----------------|
| 12.7 Your TinyTorch . . . . .                               | 141            |
| 12.8 PyTorch Lightning . . . . .                            | 142            |
| 12.9 Check Your Understanding . . . . .                     | 143            |
| 12.10Key Takeaways . . . . .                                | 145            |
| 12.11Further Reading . . . . .                              | 146            |
| 12.12What's Next . . . . .                                  | 147            |
| 12.13Get Started . . . . .                                  | 147            |
| <br><b>II Foundation Milestones</b>                         | <br><b>149</b> |
| <b>13 Historical Milestones</b>                             | <b>151</b>     |
| 13.1 Overview . . . . .                                     | 151            |
| 13.2 The Journey . . . . .                                  | 151            |
| 13.3 Why Milestones Transform Learning . . . . .            | 152            |
| 13.4 How to Use Milestones . . . . .                        | 152            |
| 13.5 Learning Philosophy . . . . .                          | 152            |
| 13.6 What's Next? . . . . .                                 | 152            |
| <br><b>14 Milestone 01: The Perceptron (1958)</b>           | <br><b>153</b> |
| 14.1 Overview . . . . .                                     | 153            |
| 14.2 What You'll Build . . . . .                            | 153            |
| 14.3 Prerequisites . . . . .                                | 153            |
| 14.4 Running the Milestone . . . . .                        | 154            |
| 14.5 Expected Results . . . . .                             | 154            |
| 14.6 The Aha Moment: Learning IS the Intelligence . . . . . | 154            |
| 14.7 Your Code Powers This . . . . .                        | 155            |
| 14.8 Historical Context . . . . .                           | 155            |
| 14.9 Systems Insights . . . . .                             | 155            |
| 14.10What's Next . . . . .                                  | 155            |
| 14.11Further Reading . . . . .                              | 155            |
| <br><b>15 Milestone 02: The XOR Crisis (1969)</b>           | <br><b>157</b> |
| 15.1 Overview . . . . .                                     | 157            |
| 15.2 What You'll Build . . . . .                            | 157            |
| 15.3 The XOR Problem . . . . .                              | 157            |
| 15.4 Prerequisites . . . . .                                | 158            |
| 15.5 Running the Milestone . . . . .                        | 158            |
| 15.6 Expected Results . . . . .                             | 158            |
| 15.7 The Aha Moment: Depth Changes Everything . . . . .     | 159            |
| 15.8 Your Code Powers This . . . . .                        | 159            |
| 15.9 Systems Insights . . . . .                             | 159            |
| 15.10Historical Context . . . . .                           | 160            |
| 15.11What's Next . . . . .                                  | 160            |
| 15.12Further Reading . . . . .                              | 160            |
| <br><b>16 Milestone 03: The MLP Revival (1986)</b>          | <br><b>161</b> |
| 16.1 Overview . . . . .                                     | 161            |
| 16.2 What You'll Build . . . . .                            | 161            |
| 16.3 Prerequisites . . . . .                                | 161            |
| 16.4 Running the Milestone . . . . .                        | 162            |
| 16.5 Expected Results . . . . .                             | 162            |
| 16.6 The Aha Moment: Automatic Feature Discovery . . . . .  | 162            |
| 16.7 Your Code Powers This . . . . .                        | 163            |
| 16.8 Historical Context . . . . .                           | 163            |

|                                         |                |
|-----------------------------------------|----------------|
| 16.9 Systems Insights . . . . .         | 163            |
| 16.10What's Next . . . . .              | 163            |
| 16.11Further Reading . . . . .          | 163            |
| <br><b>III Architecture Tier</b>        | <br><b>165</b> |
| <b>17 Module 09: Convolutions</b>       | <b>167</b>     |
| 17.1 Overview . . . . .                 | 167            |
| 17.2 Learning Objectives . . . . .      | 167            |
| 17.3 What You'll Build . . . . .        | 168            |
| 17.4 API Reference . . . . .            | 169            |
| 17.5 Core Concepts . . . . .            | 170            |
| 17.6 Common Errors . . . . .            | 175            |
| 17.7 Production Context . . . . .       | 176            |
| 17.8 Your TinyTorch . . . . .           | 176            |
| 17.9 PyTorch . . . . .                  | 177            |
| 17.10Check Your Understanding . . . . . | 178            |
| 17.11Key Takeaways . . . . .            | 180            |
| 17.12Further Reading . . . . .          | 180            |
| 17.13What's Next . . . . .              | 181            |
| 17.14Get Started . . . . .              | 182            |
| <br><b>18 Module 10: Tokenization</b>   | <br><b>183</b> |
| 18.1 Overview . . . . .                 | 183            |
| 18.2 Learning Objectives . . . . .      | 183            |
| 18.3 What You'll Build . . . . .        | 184            |
| 18.4 API Reference . . . . .            | 185            |
| 18.5 Core Concepts . . . . .            | 187            |
| 18.6 Production Context . . . . .       | 192            |
| 18.7 Your TinyTorch . . . . .           | 193            |
| 18.8 Hugging Face . . . . .             | 193            |
| 18.9 Self-Check Questions . . . . .     | 195            |
| 18.10Key Takeaways . . . . .            | 197            |
| 18.11Further Reading . . . . .          | 197            |
| 18.12What's Next . . . . .              | 198            |
| 18.13Get Started . . . . .              | 198            |
| <br><b>19 Module 11: Embeddings</b>     | <br><b>199</b> |
| 19.1 Overview . . . . .                 | 199            |
| 19.2 Learning Objectives . . . . .      | 199            |
| 19.3 What You'll Build . . . . .        | 200            |
| 19.4 API Reference . . . . .            | 201            |
| 19.5 Core Concepts . . . . .            | 202            |
| 19.6 Common Errors . . . . .            | 207            |
| 19.7 Production Context . . . . .       | 208            |
| 19.8 Your TinyTorch . . . . .           | 209            |
| 19.9 PyTorch . . . . .                  | 209            |
| 19.10Self-Check Questions . . . . .     | 211            |
| 19.11Key Takeaways . . . . .            | 212            |
| 19.12Further Reading . . . . .          | 212            |
| 19.13What's Next . . . . .              | 213            |
| 19.14Get Started . . . . .              | 214            |
| <br><b>20 Module 12: Attention</b>      | <br><b>215</b> |

|                                                          |            |
|----------------------------------------------------------|------------|
| 20.1 Overview . . . . .                                  | 215        |
| 20.2 Learning Objectives . . . . .                       | 215        |
| 20.3 What You'll Build . . . . .                         | 216        |
| 20.4 API Reference . . . . .                             | 217        |
| 20.5 Core Concepts . . . . .                             | 218        |
| 20.6 Common Errors . . . . .                             | 222        |
| 20.7 Production Context . . . . .                        | 223        |
| 20.8 Your TinyTorch . . . . .                            | 223        |
| 20.9 PyTorch . . . . .                                   | 224        |
| 20.10 Self-Check Questions . . . . .                     | 225        |
| 20.11 Key Takeaways . . . . .                            | 228        |
| 20.12 Further Reading . . . . .                          | 228        |
| 20.13 What's Next . . . . .                              | 229        |
| 20.14 Get Started . . . . .                              | 230        |
| <b>21 Module 13: Transformers</b>                        | <b>231</b> |
| 21.1 Overview . . . . .                                  | 231        |
| 21.2 Learning Objectives . . . . .                       | 231        |
| 21.3 What You'll Build . . . . .                         | 232        |
| 21.4 API Reference . . . . .                             | 233        |
| 21.5 Core Concepts . . . . .                             | 235        |
| 21.6 Production Context . . . . .                        | 240        |
| 21.7 Your TinyTorch . . . . .                            | 241        |
| 21.8 PyTorch . . . . .                                   | 241        |
| 21.9 Self-Check Questions . . . . .                      | 243        |
| 21.10 Key Takeaways . . . . .                            | 245        |
| 21.11 Further Reading . . . . .                          | 245        |
| 21.12 What's Next . . . . .                              | 246        |
| 21.13 Get Started . . . . .                              | 247        |
| <b>IV Architecture Milestones</b>                        | <b>249</b> |
| <b>22 Milestone 04: The CNN Revolution (1998)</b>        | <b>251</b> |
| 22.1 Overview . . . . .                                  | 251        |
| 22.2 What You'll Build . . . . .                         | 251        |
| 22.3 Prerequisites . . . . .                             | 252        |
| 22.4 Running the Milestone . . . . .                     | 252        |
| 22.5 Expected Results . . . . .                          | 252        |
| 22.6 The Aha Moment: Structure Matches Reality . . . . . | 253        |
| 22.7 Your Code Powers This . . . . .                     | 253        |
| 22.8 Historical Context . . . . .                        | 254        |
| 22.9 Systems Insights . . . . .                          | 254        |
| 22.10 What's Next . . . . .                              | 254        |
| 22.11 Further Reading . . . . .                          | 254        |
| <b>23 Milestone 05: The Transformer Era (2017)</b>       | <b>255</b> |
| 23.1 Overview . . . . .                                  | 255        |
| 23.2 What You'll Build . . . . .                         | 255        |
| 23.3 Prerequisites . . . . .                             | 255        |
| 23.4 Running the Milestone . . . . .                     | 256        |
| 23.5 Expected Results . . . . .                          | 256        |
| 23.6 The Aha Moment: Direct Access Everywhere . . . . .  | 256        |
| 23.7 Your Code Powers This . . . . .                     | 257        |



|           |                                      |            |
|-----------|--------------------------------------|------------|
| 23.8      | Historical Context . . . . .         | 257        |
| 23.9      | Systems Insights . . . . .           | 257        |
| 23.10     | What's Next . . . . .                | 257        |
| <b>V</b>  | <b>Optimization Tier</b>             | <b>259</b> |
| <b>24</b> | <b>Module 14: Profiling</b>          | <b>261</b> |
| 24.1      | Overview . . . . .                   | 261        |
| 24.2      | The Optimization Tier Flow . . . . . | 261        |
| 24.3      | Learning Objectives . . . . .        | 262        |
| 24.4      | What You'll Build . . . . .          | 262        |
| 24.5      | API Reference . . . . .              | 263        |
| 24.6      | Core Concepts . . . . .              | 264        |
| 24.7      | Production Context . . . . .         | 268        |
| 24.8      | Your TinyTorch . . . . .             | 268        |
| 24.9      | PyTorch . . . . .                    | 269        |
| 24.10     | Check Your Understanding . . . . .   | 270        |
| 24.11     | Key Takeaways . . . . .              | 271        |
| 24.12     | Further Reading . . . . .            | 272        |
| 24.13     | What's Next . . . . .                | 272        |
| 24.14     | Get Started . . . . .                | 273        |
| <b>25</b> | <b>Module 15: Quantization</b>       | <b>275</b> |
| 25.1      | Overview . . . . .                   | 275        |
| 25.2      | Learning Objectives . . . . .        | 275        |
| 25.3      | What You'll Build . . . . .          | 276        |
| 25.4      | API Reference . . . . .              | 277        |
| 25.5      | Core Concepts . . . . .              | 278        |
| 25.6      | Production Context . . . . .         | 282        |
| 25.7      | Your TinyTorch . . . . .             | 283        |
| 25.8      | PyTorch . . . . .                    | 283        |
| 25.9      | Check Your Understanding . . . . .   | 284        |
| 25.10     | Key Takeaways . . . . .              | 287        |
| 25.11     | Further Reading . . . . .            | 287        |
| 25.12     | What's Next . . . . .                | 288        |
| 25.13     | Get Started . . . . .                | 288        |
| <b>26</b> | <b>Module 16: Compression</b>        | <b>291</b> |
| 26.1      | Overview . . . . .                   | 291        |
| 26.2      | Learning Objectives . . . . .        | 291        |
| 26.3      | What You'll Build . . . . .          | 292        |
| 26.4      | API Reference . . . . .              | 293        |
| 26.5      | Core Concepts . . . . .              | 294        |
| 26.6      | Production Context . . . . .         | 299        |
| 26.7      | Your TinyTorch . . . . .             | 300        |
| 26.8      | PyTorch . . . . .                    | 300        |
| 26.9      | Check Your Understanding . . . . .   | 301        |
| 26.10     | Key Takeaways . . . . .              | 303        |
| 26.11     | Further Reading . . . . .            | 303        |
| 26.12     | What's Next . . . . .                | 304        |
| 26.13     | Get Started . . . . .                | 305        |
| <b>27</b> | <b>Module 17: Acceleration</b>       | <b>307</b> |
| 27.1      | Overview . . . . .                   | 307        |

|           |                                                           |            |
|-----------|-----------------------------------------------------------|------------|
| 27.2      | Why Acceleration Before Memoization?                      | 307        |
| 27.3      | Learning Objectives                                       | 308        |
| 27.4      | What You'll Build                                         | 308        |
| 27.5      | API Reference                                             | 309        |
| 27.6      | Core Concepts                                             | 310        |
| 27.7      | Common Errors                                             | 314        |
| 27.8      | Production Context                                        | 315        |
| 27.9      | Your TinyTorch                                            | 316        |
| 27.10     | PyTorch                                                   | 316        |
| 27.11     | Check Your Understanding                                  | 317        |
| 27.12     | Key Takeaways                                             | 319        |
| 27.13     | Further Reading                                           | 319        |
| 27.14     | What's Next                                               | 320        |
| 27.15     | Get Started                                               | 320        |
| <b>28</b> | <b>Module 18: Memoization</b>                             | <b>321</b> |
| 28.1      | Overview                                                  | 321        |
| 28.2      | Where Memoization Fits                                    | 321        |
| 28.3      | Learning Objectives                                       | 322        |
| 28.4      | What You'll Build                                         | 322        |
| 28.5      | API Reference                                             | 323        |
| 28.6      | Core Concepts                                             | 324        |
| 28.7      | Common Errors                                             | 328        |
| 28.8      | Production Context                                        | 329        |
| 28.9      | Your TinyTorch                                            | 329        |
| 28.10     | PyTorch                                                   | 330        |
| 28.11     | Check Your Understanding                                  | 331        |
| 28.12     | Key Takeaways                                             | 333        |
| 28.13     | Further Reading                                           | 333        |
| 28.14     | What's Next                                               | 334        |
| 28.15     | Get Started                                               | 334        |
| <b>29</b> | <b>Module 19: Benchmarking</b>                            | <b>335</b> |
| 29.1      | Overview                                                  | 335        |
| 29.2      | Learning Objectives                                       | 335        |
| 29.3      | What You'll Build                                         | 336        |
| 29.4      | API Reference                                             | 337        |
| 29.5      | Core Concepts                                             | 339        |
| 29.6      | Common Errors                                             | 343        |
| 29.7      | Production Context                                        | 345        |
| 29.8      | Your TinyTorch                                            | 345        |
| 29.9      | MLPerf (Industry Standard)                                | 346        |
| 29.10     | Check Your Understanding                                  | 347        |
| 29.11     | Key Takeaways                                             | 349        |
| 29.12     | Further Reading                                           | 349        |
| 29.13     | What's Next                                               | 350        |
| 29.14     | Get Started                                               | 351        |
| <b>VI</b> | <b>Optimization Milestones</b>                            | <b>353</b> |
| <b>30</b> | <b>Milestone 06: MLPerf — The Optimization Era (2018)</b> | <b>355</b> |
| 30.1      | Overview                                                  | 355        |
| 30.2      | What You'll Build                                         | 355        |

|                                                        |                |
|--------------------------------------------------------|----------------|
| 30.3 Prerequisites . . . . .                           | 356            |
| 30.4 Running the Milestone . . . . .                   | 356            |
| 30.5 Expected Results . . . . .                        | 356            |
| 30.6 The Aha Moment: Systematic Beats Heroic . . . . . | 357            |
| 30.7 Your Code Powers This . . . . .                   | 357            |
| 30.8 Historical Context . . . . .                      | 358            |
| 30.9 Systems Insights . . . . .                        | 358            |
| 30.10What's Next . . . . .                             | 358            |
| 30.11Further Reading . . . . .                         | 358            |
| <br><b>VII Capstone</b>                                | <br><b>359</b> |
| <b>31 Module 20: Capstone</b>                          | <b>361</b>     |
| 31.1 Overview . . . . .                                | 361            |
| 31.2 Learning Objectives . . . . .                     | 361            |
| 31.3 What You'll Build . . . . .                       | 362            |
| 31.4 API Reference . . . . .                           | 363            |
| 31.5 Core Concepts . . . . .                           | 364            |
| 31.6 Production Context . . . . .                      | 370            |
| 31.7 Your TinyTorch . . . . .                          | 371            |
| 31.8 Production MLflow . . . . .                       | 371            |
| 31.9 Check Your Understanding . . . . .                | 373            |
| 31.10Key Takeaways . . . . .                           | 375            |
| 31.11Further Reading . . . . .                         | 375            |
| 31.12What's Next . . . . .                             | 376            |
| 31.13Get Started . . . . .                             | 377            |
| <br><b>32 You Built Something Real</b>                 | <br><b>379</b> |
| 32.1 What You Accomplished . . . . .                   | 379            |
| 32.2 The Mindset Shift . . . . .                       | 379            |
| 32.3 What You Can Do Now . . . . .                     | 380            |
| 32.4 Your Code vs Production Frameworks . . . . .      | 380            |
| 32.5 Paths Forward . . . . .                           | 381            |
| 32.6 The Broader Mission . . . . .                     | 381            |
| 32.7 A Final Note . . . . .                            | 381            |
| <br><b>33 Glossary</b>                                 | <br><b>383</b> |
| 33.1 A . . . . .                                       | 383            |
| 33.2 B . . . . .                                       | 383            |
| 33.3 C . . . . .                                       | 384            |
| 33.4 D . . . . .                                       | 385            |
| 33.5 E . . . . .                                       | 385            |
| 33.6 F . . . . .                                       | 385            |
| 33.7 G . . . . .                                       | 385            |
| 33.8 H . . . . .                                       | 385            |
| 33.9 I . . . . .                                       | 385            |
| 33.10K . . . . .                                       | 386            |
| 33.11L . . . . .                                       | 386            |
| 33.12M . . . . .                                       | 386            |
| 33.13O . . . . .                                       | 387            |
| 33.14P . . . . .                                       | 387            |
| 33.15Q . . . . .                                       | 387            |
| 33.16R . . . . .                                       | 387            |

|                  |     |
|------------------|-----|
| 33.17S . . . . . | 388 |
| 33.18T . . . . . | 388 |
| 33.19V . . . . . | 388 |

# List of Figures

---

|      |                                                                                                                                                                                                                                                            |     |
|------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----|
| 3.1  | <b>TinyTorch Module Flow.</b> The 20 modules progress through three tiers: Foundation (blue) builds core ML primitives, Architecture (green) applies them to vision and language tasks, and Optimization (orange) makes systems production-ready. . . . .  | 8   |
| 4.1  | <b>Your TinyTorch Journey:</b> Install once, then loop through start → complete → milestone for each of the 20 modules — every milestone runs on the code you just wrote. . . . .                                                                          | 11  |
| 5.1  | <b>Tensor Class Architecture:</b> One class packages data, shape metadata, and every operation the rest of TinyTorch will lean on. . . . .                                                                                                                 | 20  |
| 5.2  | <b>Broadcasting Mechanics:</b> How a vector is expanded to match a matrix's shape during arithmetic operations. . . . .                                                                                                                                    | 23  |
| 5.3  | <b>The Tensor Stack:</b> From the Python interface down to the hardware CPU instructions. . . . .                                                                                                                                                          | 28  |
| 6.1  | <b>Activation functions in TinyTorch:</b> ReLU, Sigmoid, Tanh, GELU, and Softmax transformations. . . . .                                                                                                                                                  | 36  |
| 7.1  | <b>One interface, three subclasses.</b> Linear, Dropout, and Sequential all inherit <code>forward()</code> and <code>parameters()</code> from a single <code>Layer</code> base — Sequential adds composition by chaining the others into networks. . . . . | 52  |
| 8.1  | <b>TinyTorch Loss Functions:</b> MSE for regression, and Cross-Entropy for classification tasks. . . . .                                                                                                                                                   | 70  |
| 9.1  | <b>TinyTorch Data Pipeline:</b> From raw dataset storage to training-ready batches. . . . .                                                                                                                                                                | 86  |
| 10.1 | <b>TinyTorch Autograd Engine:</b> Reverse-mode automatic differentiation infrastructure. . . . .                                                                                                                                                           | 102 |
| 11.1 | <b>TinyTorch Optimizer Hierarchy:</b> From basic SGD to advanced adaptive algorithms. . . . .                                                                                                                                                              | 118 |
| 12.1 | <b>TinyTorch Training Ecosystem:</b> Orchestration of scheduling, clipping, and checkpointing within the Trainer class. . . . .                                                                                                                            | 132 |
| 17.1 | <b>TinyTorch Spatial Operations:</b> Convolutional and pooling layers for hierarchical visual feature extraction. . . . .                                                                                                                                  | 168 |
| 18.1 | <b>TinyTorch Tokenization Infrastructure:</b> Converting raw text into model-ready numerical sequences. . . . .                                                                                                                                            | 184 |
| 19.1 | <b>TinyTorch Embedding System:</b> Mapping tokens and positions to learned continuous representations. . . . .                                                                                                                                             | 200 |
| 20.1 | <b>Attention Mechanism:</b> Information retrieval paradigm where similarity scores between Queries and Keys determine how much of each Value is retrieved. . . . .                                                                                         | 216 |
| 21.1 | . . . . .                                                                                                                                                                                                                                                  | 232 |
| 24.1 | <b>TinyTorch Profiling System:</b> Tools for measuring execution time and memory allocation. . . . .                                                                                                                                                       | 262 |
| 25.1 | <b>TinyTorch Quantization System:</b> Methods for converting models to lower precision. . . . .                                                                                                                                                            | 276 |

|      |                                                                                                                                                                                                                    |     |
|------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----|
| 26.1 | <b>TinyTorch Compression System:</b> Reducing model size via pruning and distillation. . . . .                                                                                                                     | 292 |
| 27.1 | <b>TinyTorch Acceleration System:</b> Vectorized kernels and SIMD optimization. . . . .                                                                                                                            | 308 |
| 28.1 | <b>TinyTorch KV Caching:</b> Reusing previous computations to speed up text generation. . . . .                                                                                                                    | 322 |
| 29.1 | <b>TinyTorch Benchmarking System:</b> Statistical measurement of performance and accuracy. . .                                                                                                                     | 336 |
| 31.1 | <b>TinyTorch Submission Pipeline:</b> Baseline and optimized models flow through <code>BenchmarkReport</code> and <code>generate_submission()</code> to produce a schema-validated <code>results.json</code> . . . | 362 |

# List of Tables

---

|      |                                                                                        |     |
|------|----------------------------------------------------------------------------------------|-----|
| 3.1  | Concrete outcomes unlocked at each module checkpoint. . . . .                          | 9   |
| 4.1  | Milestones unlocked as learners complete more modules. . . . .                         | 14  |
| 4.2  | TinyTorch module tiers with scope and time estimate. . . . .                           | 15  |
| 5.1  | Five-part implementation roadmap for the Tensor class. . . . .                         | 20  |
| 5.2  | Read-only properties exposed by the Tensor class. . . . .                              | 21  |
| 5.3  | Arithmetic dunder methods and their Python operator equivalents. . . . .               | 21  |
| 5.4  | Matrix and shape manipulation methods on Tensor. . . . .                               | 22  |
| 5.5  | Reduction methods that collapse tensor dimensions. . . . .                             | 22  |
| 5.6  | Tensor ranks from scalar to 4D with concrete examples. . . . .                         | 22  |
| 5.7  | Broadcasting compatibility for representative shape pairs. . . . .                     | 24  |
| 5.8  | Memory and time cost of view operations versus copying ops. . . . .                    | 24  |
| 5.9  | Computational complexity of core tensor operations. . . . .                            | 27  |
| 5.10 | Feature comparison between TinyTorch Tensor and PyTorch Tensor. . . . .                | 30  |
| 6.1  | Implementation roadmap for the core activation functions. . . . .                      | 36  |
| 6.2  | Mathematical form, output range, and use case for each activation. . . . .             | 37  |
| 6.3  | Computational cost of activations relative to ReLU. . . . .                            | 43  |
| 6.4  | Feature comparison between TinyTorch activations and PyTorch equivalents. . . . .      | 43  |
| 6.5  | How activations feed into subsequent TinyTorch modules. . . . .                        | 49  |
| 7.1  | Implementation roadmap for the Layer, Linear, Dropout, and Sequential classes. . . . . | 52  |
| 7.2  | Methods defined by the Layer base class. . . . .                                       | 53  |
| 7.3  | Methods on the Linear layer. . . . .                                                   | 54  |
| 7.4  | Methods on the Dropout layer. . . . .                                                  | 54  |
| 7.5  | Methods on the Sequential container. . . . .                                           | 54  |
| 7.6  | Compute and memory cost of Linear and Dropout forward passes. . . . .                  | 59  |
| 7.7  | Feature comparison between TinyTorch layers and PyTorch's nn module. . . . .           | 61  |
| 7.8  | How the Layer abstractions feed into later training modules. . . . .                   | 66  |
| 8.1  | Implementation roadmap for the three loss functions. . . . .                           | 70  |
| 8.2  | Constructor, forward signature, and use case for each loss. . . . .                    | 71  |
| 8.3  | Input and output shape contracts for each loss function. . . . .                       | 72  |
| 8.4  | Feature comparison between TinyTorch losses and PyTorch equivalents. . . . .           | 76  |
| 8.5  | How losses feed into subsequent training modules. . . . .                              | 83  |
| 9.1  | Implementation roadmap for the Dataset and DataLoader classes. . . . .                 | 86  |
| 9.2  | Required methods on the Dataset abstract base class. . . . .                           | 87  |
| 9.3  | Core methods on the DataLoader class. . . . .                                          | 88  |
| 9.4  | Feature comparison between TinyTorch DataLoader and PyTorch DataLoader. . . . .        | 93  |
| 9.5  | How the DataLoader feeds into subsequent training modules. . . . .                     | 99  |
| 10.1 | Implementation roadmap for the reverse-mode autograd engine. . . . .                   | 102 |
| 10.2 | Backward Function classes and their gradient rules. . . . .                            | 103 |
| 10.3 | Methods added to the Tensor class for autograd. . . . .                                | 104 |

|      |                                                                                       |     |
|------|---------------------------------------------------------------------------------------|-----|
| 10.4 | Global helper functions for enabling autograd. . . . .                                | 104 |
| 10.5 | Feature comparison between TinyTorch autograd and PyTorch autograd. . . . .           | 109 |
| 10.6 | How autograd feeds into subsequent optimizer and training modules. . . . .            | 115 |
| 11.1 | Implementation roadmap for the SGD and Adam optimizers. . . . .                       | 118 |
| 11.2 | Methods defined by the Optimizer base class. . . . .                                  | 119 |
| 11.3 | State-management methods on the SGD optimizer. . . . .                                | 119 |
| 11.4 | Feature comparison between TinyTorch optimizers and torch.optim. . . . .              | 124 |
| 11.5 | How optimizers feed into subsequent training modules. . . . .                         | 129 |
| 12.1 | Implementation roadmap for the Trainer and learning-rate schedule. . . . .            | 132 |
| 12.2 | Method on the CosineSchedule class. . . . .                                           | 133 |
| 12.3 | Core methods on the Trainer class. . . . .                                            | 133 |
| 12.4 | Memory footprint of model parameters, gradients, and optimizer state. . . . .         | 140 |
| 12.5 | Feature comparison between TinyTorch Trainer and PyTorch training stacks. . . . .     | 141 |
| 12.6 | How the Trainer gets reused in the Architecture tier modules. . . . .                 | 147 |
| 13.1 | Historical milestone timeline and required modules for each. . . . .                  | 151 |
| 14.1 | Prerequisite modules for the Perceptron milestone. . . . .                            | 154 |
| 14.2 | Expected accuracy for the Perceptron milestone scripts. . . . .                       | 154 |
| 14.3 | TinyTorch components that power the Perceptron milestone. . . . .                     | 155 |
| 15.1 | Prerequisite modules for the XOR milestone. . . . .                                   | 158 |
| 15.2 | Expected loss and accuracy for the XOR milestone scripts. . . . .                     | 159 |
| 15.3 | TinyTorch components that power the XOR milestone. . . . .                            | 159 |
| 16.1 | Prerequisite modules for the MLP milestone. . . . .                                   | 162 |
| 16.2 | Expected accuracy and training time for the MLP milestone scripts. . . . .            | 162 |
| 16.3 | TinyTorch components that power the MLP milestone. . . . .                            | 163 |
| 17.1 | Implementation roadmap for Conv2d, MaxPool2d, and supporting layers. . . . .          | 168 |
| 17.2 | Core methods on spatial layers (Conv2d, MaxPool2d, BatchNorm2d). . . . .              | 170 |
| 17.3 | Computational complexity of 2D spatial operations. . . . .                            | 174 |
| 17.4 | Feature comparison between TinyTorch Conv2d and PyTorch cuDNN. . . . .                | 176 |
| 17.5 | How spatial ops feed into later milestone and optimization modules. . . . .           | 181 |
| 18.1 | Implementation roadmap for the Tokenizer classes. . . . .                             | 184 |
| 18.2 | Methods on the CharTokenizer class. . . . .                                           | 185 |
| 18.3 | Methods on the BPETokenizer class. . . . .                                            | 186 |
| 18.4 | Internal helper methods used by the BPETokenizer. . . . .                             | 186 |
| 18.5 | Utility functions for tokenizer creation and analysis. . . . .                        | 187 |
| 18.6 | Complexity and speed of character and BPE tokenization stages. . . . .                | 191 |
| 18.7 | Vocabulary size and tokenizer strategy in production LLMs. . . . .                    | 192 |
| 18.8 | Feature comparison between TinyTorch tokenizers and Hugging Face Tokenizers. . . . .  | 192 |
| 18.9 | How tokenization feeds into embeddings, attention, and transformers. . . . .          | 198 |
| 19.1 | Implementation roadmap for the embedding classes. . . . .                             | 200 |
| 19.2 | Core methods on the Embedding class. . . . .                                          | 201 |
| 19.3 | Core methods on the PositionalEncoding class. . . . .                                 | 201 |
| 19.4 | Core methods on the combined token+position Embedding wrapper. . . . .                | 202 |
| 19.5 | Vocabulary size, embedding dimension, and memory cost across production LLMs. . . . . | 206 |
| 19.6 | Feature comparison between TinyTorch embeddings and torch.nn.Embedding. . . . .       | 208 |
| 19.7 | How embeddings feed into attention and transformer modules. . . . .                   | 214 |
| 20.1 | Implementation roadmap for scaled dot-product and multi-head attention. . . . .       | 216 |



|      |                                                                                         |     |
|------|-----------------------------------------------------------------------------------------|-----|
| 20.2 | Core methods on the MultiHeadAttention class. . . . .                                   | 217 |
| 20.3 | Time and memory complexity of the three attention computation steps. . . . .            | 221 |
| 20.4 | Feature comparison between TinyTorch attention and nn.MultiheadAttention. . . . .       | 223 |
| 20.5 | How attention feeds into the transformer block assembly. . . . .                        | 229 |
| 21.1 | Implementation roadmap for LayerNorm, MLP, TransformerBlock, and GPT. . . . .           | 232 |
| 21.2 | Core methods on the LayerNorm class. . . . .                                            | 234 |
| 21.3 | Core methods on the MLP feed-forward block. . . . .                                     | 234 |
| 21.4 | Core methods on the TransformerBlock class. . . . .                                     | 234 |
| 21.5 | Core methods on the complete GPT model class. . . . .                                   | 235 |
| 21.6 | Parameter count breakdown for a single 512-dim, 8-head transformer block. . . . .       | 239 |
| 21.7 | Attention memory growth with sequence length for a fixed-size transformer. . . . .      | 240 |
| 21.8 | Feature comparison between TinyTorch transformers and production PyTorch. . . . .       | 241 |
| 21.9 | How transformers feed into profiling, quantization, and capstone modules. . . . .       | 246 |
| 22.1 | Prerequisite modules for the CNN milestone. . . . .                                     | 252 |
| 22.2 | Expected accuracy for the CNN milestone on TinyDigits and CIFAR-10. . . . .             | 252 |
| 22.3 | TinyTorch components that power the CNN milestone. . . . .                              | 253 |
| 23.1 | Prerequisite modules for the Transformer milestone. . . . .                             | 256 |
| 23.2 | Expected success criteria and runtime for the Transformer milestone. . . . .            | 256 |
| 23.3 | TinyTorch components that power the Transformer milestone. . . . .                      | 257 |
| 24.1 | Implementation roadmap for the Profiler class and its measurement methods. . . . .      | 262 |
| 24.2 | Core measurement methods on the Profiler class. . . . .                                 | 263 |
| 24.3 | Higher-level analysis methods on the Profiler class. . . . .                            | 264 |
| 24.4 | Utility functions for quick profiling and weight analysis. . . . .                      | 264 |
| 24.5 | Feature comparison between TinyTorch Profiler and PyTorch profiling tools. . . . .      | 268 |
| 24.6 | How the profiler feeds into optimization-tier modules. . . . .                          | 272 |
| 25.1 | Implementation roadmap for INT8 quantization and QuantizedLinear. . . . .               | 276 |
| 25.2 | Core methods on the QuantizedLinear class. . . . .                                      | 277 |
| 25.3 | Model-level quantization helper functions. . . . .                                      | 278 |
| 25.4 | Core methods on the Quantizer convenience class. . . . .                                | 278 |
| 25.5 | Feature comparison between TinyTorch quantizer and PyTorch Quantization. . . . .        | 282 |
| 25.6 | How quantization stacks with compression, acceleration, and capstone modules. . . . .   | 288 |
| 26.1 | Implementation roadmap for pruning and knowledge distillation. . . . .                  | 292 |
| 26.2 | Magnitude and structured pruning functions. . . . .                                     | 293 |
| 26.3 | Key method on the knowledge-distillation helper class. . . . .                          | 293 |
| 26.4 | Compression, accuracy, and speedup trade-offs across compression techniques. . . . .    | 299 |
| 26.5 | Feature comparison between TinyTorch compression and PyTorch pruning utilities. . . . . | 299 |
| 26.6 | How compression feeds into acceleration, memoization, and benchmarking. . . . .         | 304 |
| 27.1 | Implementation roadmap for vectorized kernels and operator fusion. . . . .              | 308 |
| 27.2 | Fused and baseline GELU kernel functions. . . . .                                       | 309 |
| 27.3 | Cache-aware matrix multiplication function. . . . .                                     | 309 |
| 27.4 | Arithmetic intensity and optimization strategy by operation. . . . .                    | 314 |
| 27.5 | Feature comparison between TinyTorch acceleration and PyTorch internals. . . . .        | 315 |
| 27.6 | How acceleration composes with memoization, benchmarking, and capstone. . . . .         | 320 |
| 28.1 | Implementation roadmap for the KV cache data structure. . . . .                         | 322 |
| 28.2 | Core methods on the KVCache class. . . . .                                              | 323 |
| 28.3 | Helper functions for enabling and disabling the KV cache. . . . .                       | 324 |
| 28.4 | KV-cache memory and compute savings as sequence length grows. . . . .                   | 327 |

|      |                                                                                      |     |
|------|--------------------------------------------------------------------------------------|-----|
| 28.5 | Feature comparison between TinyTorch KV cache and PyTorch transformers. . . . .      | 329 |
| 28.6 | How memoization stacks with quantization, acceleration, and benchmarking. . . . .    | 334 |
| 29.1 | Implementation roadmap for the benchmarking suite. . . . .                           | 336 |
| 29.2 | Statistical properties computed by BenchmarkResult. . . . .                          | 337 |
| 29.3 | Serialization methods on the BenchmarkResult dataclass. . . . .                      | 337 |
| 29.4 | Core measurement methods on the Benchmark class. . . . .                             | 338 |
| 29.5 | Methods on the BenchmarkSuite class for multi-metric evaluation. . . . .             | 339 |
| 29.6 | Feature comparison between TinyTorch benchmarking and MLPerf practices. . . . .      | 345 |
| 29.7 | How benchmarking powers each capstone competition event. . . . .                     | 350 |
| 30.1 | Prerequisite modules for the MLPerf milestone. . . . .                               | 356 |
| 30.2 | Expected size and accuracy at each static optimization stage. . . . .                | 356 |
| 30.3 | Per-token generation speed with and without the KV cache. . . . .                    | 357 |
| 30.4 | TinyTorch components that power the MLPerf milestone. . . . .                        | 357 |
| 31.1 | Implementation roadmap for the capstone benchmarking artifact. . . . .               | 362 |
| 31.2 | Data properties stored on the BenchmarkReport class. . . . .                         | 363 |
| 31.3 | Core methods on the BenchmarkReport class. . . . .                                   | 364 |
| 31.4 | Feature comparison between TinyTorch capstone benchmarking and industry systems. . . | 371 |
| 31.5 | Suggested directions for extending the framework beyond the capstone. . . . .        | 376 |
| 32.1 | Component-by-component comparison of TinyTorch and PyTorch internals. . . . .        | 380 |

## Chapter 1

# About This Guide

---

**TinyTorch** is a build-it-yourself companion to the *Machine Learning Systems* textbook (available at [mlsysbook.ai](https://mlsysbook.ai)). This guide collects the orientation, quick-start, and 20 implementation modules into a single PDF for offline reading and print study.

The modules progress through three tiers — **Foundation** (01-08) teaches the primitives (tensor, autograd, optimizer, training loop), **Architecture** (09-13) applies them to vision and language, and **Optimization** (14-19) makes the resulting systems fast, small, and measured. The **Capstone** (Module 20) ties everything back into a deployable framework.

This PDF is auto-generated from the live site at [mlsysbook.ai/tinytorch/](https://mlsysbook.ai/tinytorch/). The web version is authoritative — it includes interactive widgets, embedded slides, and quizzes that are omitted here. Use this guide when you need a portable reference; use the site when you're actually working through the modules.

---

© 2024-2026 Harvard University. Licensed under [CC-BY-NC-SA 4.0](https://creativecommons.org/licenses/by-nc-sa/4.0/).



## Chapter 2

# Welcome

---

Everyone wants to be an astronaut. Very few want to be the rocket scientist.

Machine learning is no different. Everyone wants to train models, run inference, deploy AI. Few want to understand how the frameworks actually work. Fewer still want to build one.

The world has plenty of users. It does not have enough builders: people who can debug, optimize, and adapt systems when the black box breaks down.

TinyTorch is for the builders.

## 2.1 The Problem

Most people can use PyTorch or TensorFlow. They can import libraries, call functions, train models. But very few understand how these frameworks work: how memory is managed for tensors, how autograd builds computation graphs, how optimizers update parameters. And almost no one has a guided, structured way to learn that from the ground up.

Users hit walls that builders do not:

- Out of memory? You need to understand tensor allocation.
- Gradients exploding? You need to understand the computation graph.
- Training too slow? You need to find the bottleneck.
- Deploying on a microcontroller? You need to know what can be stripped away.

The framework becomes a black box you cannot debug, optimize, or adapt. You are stuck waiting for someone else to solve your problem.

Students cannot learn this from production code. PyTorch is too large, too complex, too optimized. Fifty thousand lines of C++ across hundreds of files. No one learns to build rockets by studying the Saturn V.

They also cannot learn it from toy scripts. A hundred-line neural network does not reveal the architecture of a framework. It hides it.

## 2.2 The Solution: AI Bricks

TinyTorch teaches you the **AI bricks**: the stable engineering foundations you can use to build any AI system. Small enough to learn from: bite-sized code that runs even on a Raspberry Pi. Big enough to matter: showing the real architecture of how frameworks are built.

**MLSysBook.** The **Machine Learning Systems** textbook teaches the *concepts* of the rocket ship: propulsion, guidance, life support.

**TinyTorch.** This is where you actually *build* a small rocket with your own hands. Not a toy. A real framework.

This is how you move from *using* machine learning to *engineering* it: from running code in a notebook to designing the systems that run underneath.

## 2.3 Who This Is For

**Students & Researchers.** Want to understand ML systems deeply, not just use them superficially. If you have wondered “how does that actually work?”, this is for you.

**ML Engineers.** Need to debug, optimize, and deploy models in production. Understanding the systems underneath makes you more effective.

**Systems Programmers.** You understand memory hierarchies, computational complexity, performance optimization. You want to apply it to ML.

**Self-taught Engineers.** Can use frameworks but want to know how they work. Preparing for ML infrastructure roles and need systems-level understanding.

What you need is not another API tutorial. You need to build.

## 2.4 What You Need to Start

You do not need to be a machine learning expert, and you do not need to have built a framework before. You need three things.

**Python.** If you can read and write a function, a loop, and a class, you have enough. There is no C++, no CUDA, no assembly. Everything sits on NumPy, and you will pick up the NumPy you need as you go.

**The math you already have.** Vectors, matrices, and the chain rule from first-year calculus. That is the whole list. No measure theory, no convex optimization, no proofs. When a derivative shows up, we derive it in code, not on a chalkboard.

**A machine that turns on.** TinyTorch runs on a laptop. It runs on a Raspberry Pi. The code is small on purpose, so you can hold a module in your head and run it on hardware you already own. No GPU, no cloud account, no four-figure compute bill.

If you have used PyTorch and felt like a tourist, this is the ground you were standing on.

## 2.5 What You Will Build

By the end of TinyTorch, you will have implemented:

- A tensor library with broadcasting, reshaping, and matrix operations
- Activation functions with numerical stability considerations
- Neural network layers: linear, convolutional, normalization
- An autograd engine that builds computation graphs and computes gradients
- Optimizers that update parameters using those gradients
- Data loaders that handle batching, shuffling, and preprocessing
- A complete training loop that ties everything together
- Tokenizers, embeddings, attention, and transformer architectures
- Profiling, quantization, and optimization techniques

Not a simulation. The actual architecture of modern ML frameworks, implemented at a scale you can hold in your head.

## 2.6 How to Learn

Each module follows a **Build-Use-Reflect** cycle: implement from scratch, apply to real problems, then connect what you built to production systems and understand the tradeoffs. Work through Foundation first, then choose your path based on your interests.

**Type every line yourself.** Do not copy-paste. The learning happens in the struggle of implementation.

**Profile your code.** Use the built-in profiling tools. Measure first, optimize second.

**Run the tests.** Every module ships with tests. When they pass, you have built something real.

**Compare with PyTorch.** Once your implementation works, compare with PyTorch’s equivalent to see how production frameworks scale the same ideas.

Take your time. The goal is not to finish fast. The goal is to understand deeply.

*Building systems creates irreversible understanding.*

## 2.7 The Bigger Picture

TinyTorch is one piece of a larger curriculum, and every piece exists for the same reason: students who only read do not internalize, and students who only code do not generalize. The [Machine Learning Systems](#) textbook gives you the concepts: how training works, why accelerators matter, what makes inference cheap or expensive. TinyTorch makes you build the machinery yourself. The hardware kits put what you built on real devices, where memory limits, power budgets, and latency stop being abstractions. And StaffML tests whether you can reason about these systems under pressure, the way an interview or a production incident will.

This follows a long tradition in engineering education. You learned electronics by wiring a circuit on a breadboard. You learned architecture by laying out a processor on an FPGA. You learned operating systems by building a small kernel of your own, your first TinyOS. You do not understand a system until you have built one. The same tradition runs through systems education: SICP’s “build to understand” philosophy, xv6’s transparent operating system, Nachos, Pintos. TinyTorch brings it to machine learning, and it grew out of years of teaching these ideas at Harvard and building the open MLSysBook curriculum. The pedagogical principles are detailed in our [research paper](#), which positions this work within decades of CS education research.

The next generation of engineers cannot rely on magic. They need to see how everything fits together, from a single tensor allocation up to a full training loop, and feel that the systems running modern AI are not an unreachable tower but something they can open, shape, and rebuild.

That is what TinyTorch offers: the confidence that comes from having built it yourself.

## 2.8 For Instructors

Every engineering course has a lab, because students learn to build by building. TinyTorch is that lab for machine learning systems, and it is built to be taught. The modules come with tests and autograding so you can guide students as they build, and the instructor hub at [mlsysbook.ai](https://mlsysbook.ai) collects what you need to bring this into a classroom: the AI Engineering Blueprint, lecture slides, a course map, and assessment guides. These materials are still growing, and we welcome ideas and contributions from anyone teaching with them.

## 2.9 Make It Better

TinyTorch is open source, and it is built the way it teaches: in the open, by people who wanted to understand it. Every module, test, and milestone lives in the open, so if you find a bug, an explanation that did not land, or a cleaner way to teach a concept, you can fix it, and the next reader gets the benefit. The students who learn the most from TinyTorch are often the ones who end up improving it.

Found a problem? Built something better? Think a design choice is wrong? Bring it to [mlsysbook.ai/git](https://mlsysbook.ai/git).

*Prof. Vijay Janapa Reddi (Harvard University) 2025*

## 2.10 What’s Next?

**See the Big Picture** →: how all 20 modules connect, what you will build, and which path to take.





## Chapter 3

# Big Picture

---

### 2-minute orientation before you begin building

You're about to build a working ML framework, one module at a time. Before diving in, take two minutes to see how the twenty modules connect, what you'll have when you're done, and which path through the book fits your goals.

### 3.1 The Journey: Foundation to Production

TinyTorch takes you from a bare tensor to a production-style ML system in twenty modules. They connect like this.

#### Three tiers, one system:

- **Foundation (01-08)** — Build the core machinery. Tensors hold data, activations add non-linearity, layers combine them, losses measure error, DataLoader streams batches, autograd computes gradients, optimizers update weights, training orchestrates the loop.
- **Architecture (green, 09-13)** — Apply the foundation to real problems. The DataLoader from Module 05 feeds data; from there you take one of two paths—convolutions for images, or the transformer stack (Tokenization → Embeddings → Attention → Transformers) for text.
- **Optimization (14-19)** — Make it fast. Profile to find bottlenecks, then apply quantization, compression, acceleration, or memoization. Benchmark to prove the gain.

Figure 3.1 shows how the pieces fit together.




Figure 3.1: **TinyTorch Module Flow.** The 20 modules progress through three tiers: Foundation (blue) builds core ML primitives, Architecture (green) applies them to vision and language tasks, and Optimization (orange) makes systems production-ready.

#### Flexible paths:

- **Vision focus** — Foundation (01-08) → Convolutions (09) → Optimization (14-19)
- **Language focus** — Foundation (01-08) → Tokenization → Embeddings → Attention → Transformers (10-13) → Optimization (14-19)
- **Full course** — Both paths → Capstone (20)

## 3.2 Milestones You'll Unlock

As you build, you unlock historical milestones—moments when your code does something that once made headlines:

1. **1958 Perceptron:** Your first learning algorithm with automatic weight updates (Rosenblatt)
2. **1969 XOR:** Your MLP solves the problem that stumped single-layer networks (Minsky & Papert → Rumelhart)
3. **1986 MLP:** Your network recognizes handwritten digits on real data
4. **1998 CNN:** Your convolutional network classifies images with spatial understanding (LeCun's LeNet-5)
5. **2017 Transformer:** Your attention mechanism solves structured sequence tasks (Vaswani et al.)
6. **2018 MLPerf:** Your optimized system benchmarks at production speed

Each milestone activates when you complete the required modules. You're not just learning—you're recreating seventy years of ML evolution, one working implementation at a time.

## 3.3 What You'll Have at the End

Concrete outcomes at each major checkpoint:

Table 3.1 pins down the concrete outcome you unlock at each checkpoint.

Table 3.1: Concrete outcomes unlocked at each module checkpoint.

| After Module  | You'll Have Built                            | Historical Context                 |
|---------------|----------------------------------------------|------------------------------------|
| 01-03         | Perceptron forward pass                      | Rosenblatt 1958                    |
| 01-08         | MLP solving XOR + complete training pipeline | AI Winter breakthrough 1969→1986   |
| 01-09         | CNN with convolutions and pooling            | LeNet-5 (1998)                     |
| 01-08 + 11-13 | GPT model with autoregressive generation     | "Attention Is All You Need" (2017) |
| 01-08 + 14-19 | Optimized, quantized, accelerated system     | Production ML today                |
| 01-20         | MLPerf-style benchmarking submission         | Torch Olympics                     |

#### 💡 The North Star Build

By module 13, you'll have a complete GPT model generating text—built from raw Python. By module 20, you'll benchmark your entire framework with MLPerf-style submissions. Every tensor operation, every gradient calculation, every optimization trick: **you wrote it**.

## 3.4 Choose Your Learning Path

Pick the route that matches your goals and available time.

**Sequential Builder** — Complete all 20 modules in order. *Best for:* students, career transitioners, deep understanding. *Time:* 60–80 hours (8–12 weeks part-time). *Outcome:* a complete mental model of ML systems.

**Vision Track** — Modules 01–09 then 14–19 (CNNs + optimization). *Best for:* computer vision focus, MLOps practitioners. *Time:* 40–50 hours. *Outcome:* CNN architectures with production optimization.

**Language Track** — Modules 01–08 then 10–13 (transformers + GPT). *Best for:* NLP focus, research engineers. *Time:* 35–45 hours. *Outcome:* a complete GPT model with text generation.

**Instructor Sampler** — Read modules 01, 03, 05, 07, 12 (key concepts). *Best for:* evaluating for course adoption. *Time:* 8–12 hours (reading, not building). *Outcome:* assessment of the pedagogical approach.

#### 💡 All paths start at Module 01

Module 01 (Tensor) is the foundation everything else builds on. Start there, then switch paths anytime based on what you find interesting.

## 3.5 Expect to Struggle (That's the Design)

#### ! Getting stuck is not a bug—it's a feature

TinyTorch treats productive struggle as a teaching tool. You will debug tensor shape mismatches, trace gradient flow through tangled graphs, and fight for memory inside tight constraints. The friction is intentional. It is your brain rewiring around how ML systems actually work.

**What helps when you're stuck:**

- Run the tests early and often—they're your fastest feedback loop.
- The `if __name__ == "__main__":` blocks show the expected workflow.
- The ML Systems Thinking questions validate that you understood, not just that you typed.
- Production context notes connect your implementation back to PyTorch and TensorFlow.

**When to ask for help:**

- After you've run the tests and read the error message carefully.
- After you've tried explaining the problem out loud to a rubber duck.
- If you've been stuck on a single bug for more than thirty minutes.

The goal isn't to never struggle. It's to struggle *productively*, and to leave each module knowing why the working version works.

### 3.6 Start Building

You have the map. Module 01 builds the tensor—the data structure every other module depends on. A few hours from now you'll have a working `Tensor` class and a green test suite, and the path to CNNs, transformers, and an MLPerf-style benchmark will be one module shorter.

**Next step.** Follow the [Quick Start Guide](#) to set up your environment (2 minutes), complete Module 01: Tensor (2–3 hours), and watch your first tests pass.

**i Before you start**

You don't need to be an expert. You need to be curious and willing to struggle through hard problems. If you want to know *why* the book is built this way before you write a line of code, read the [Learning Philosophy](#) first.

The journey from tensors to transformers starts with a single `import tinytorch`.

## 🔥 Chapter 4

# Getting Started with TinyTorch

### 📘 Prerequisites Check

This guide requires **Python programming** (classes, functions, NumPy basics) and **basic linear algebra** (matrix multiplication).

## 4.1 The Journey

TinyTorch follows a simple pattern: **build modules, unlock milestones, recreate ML history**.




Figure 4.1: **Your TinyTorch Journey**: Install once, then loop through start → complete → milestone for each of the 20 modules — every milestone runs on the code you just wrote.

As you complete modules, you unlock milestones that recreate landmark moments in ML history—using YOUR code.

## 4.2 Step 1: Install & Setup (2 Minutes)

### 4.3 macOS / Linux

```
# Install TinyTorch (run from a project folder like ~/projects)
curl -sSL mlsysbook.ai/tinytorch/install.sh | bash

# Activate and verify
cd tinytorch
source .venv/bin/activate
tito setup
```

## 4.4 Windows

TinyTorch works on Windows using **Git Bash** (included with Git for Windows).

**Step 1: Install Git for Windows** (if you don't have it) - Download from [git-scm.com/download/win](https://git-scm.com/download/win) - Run the installer with default options

**Step 2: Open Git Bash**

- Search “Git Bash” in the Start menu and open it

**Step 3: Install TinyTorch**

```
# In Git Bash (run from a project folder like ~/projects)
curl -sSL mlsysbook.ai/tinytorch/install.sh | bash

# Activate and verify
cd tinytorch
source .venv/Scripts/activate
tito setup
```

**What this does:**

- Checks your system (Python 3.10+, git)
- Downloads TinyTorch to a `tinytorch/` folder
- Creates an isolated virtual environment
- Installs all dependencies
- Verifies installation

**Check your version:**

```
tito --version
```

**Update TinyTorch:**

```
tito system update
```

## 4.5 Step 2: Your First Module (15 Minutes)

Let's build Module 01 (Tensor)—the foundation of all neural networks.

### 4.5.1 Start the module

```
tito module start 01
```

This opens the module notebook and tracks your progress.

### 4.5.2 Work in the notebook

Edit `modules/01_tensor/tensor.ipynb` in Jupyter:

```
jupyter lab modules/01_tensor/tensor.ipynb
```

You'll implement: - N-dimensional array creation - Mathematical operations (add, multiply, matmul) - Shape manipulation (reshape, transpose)

### 4.5.3 Complete the module

When your implementation is ready, export it to the TinyTorch package:

```
tito module complete 01
```

Your code is now importable:

```
from tinytorch.core.tensor import Tensor  # YOUR implementation!
x = Tensor([1, 2, 3])
```

## 4.6 Step 3: Your First Milestone

Now for the payoff! After completing the required modules (01-03), run a milestone:

```
tito milestone run perceptron
```

The milestone uses YOUR implementations to recreate Rosenblatt's 1958 Perceptron:

```
Checking prerequisites for Milestone 01...
All required modules completed!

Testing YOUR implementations...
  * Tensor import successful
  * Activations import successful
  * Layers import successful
YOUR TinyTorch is ready!

+----- Milestone 01 (1958) -----+
| Milestone 01: Perceptron (1958) |
| Frank Rosenblatt's First Neural Network |
| |
| Running: milestones/01_1958_perceptron/01_rosenblatt_forward.py |
| All code uses YOUR TinyTorch implementations! |
+-----+

Starting Milestone 01...

Assembling perceptron with YOUR TinyTorch modules...
  * Linear layer: 2 -> 1 (YOUR Module 03!)
  * Activation: Sigmoid (YOUR Module 02!)

+----- Achievement Unlocked -----+
| MILESTONE ACHIEVED! |
| |
| You completed Milestone 01: Perceptron (1958) |
| Frank Rosenblatt's First Neural Network |
| |
| What makes this special: |
```

```
| - Every tensor operation: YOUR Tensor class |
| - Every layer: YOUR Linear implementation |
| - Every activation: YOUR Sigmoid function |
+-----+
```

You're recreating ML history with your own code. *By Module 19, you'll benchmark against MLPerf—the industry standard for ML performance.*

## 4.7 The Pattern Continues

As you complete more modules, you unlock more milestones:

Table 4.1 shows which milestones unlock as you complete more modules.

Table 4.1: **Milestones unlocked as learners complete more modules.**

| Modules Completed | Milestone             | What You Recreate                            |
|-------------------|-----------------------|----------------------------------------------|
| 01-03             | Perceptron (1958)     | First neural network (forward pass)          |
| 01-03             | XOR Crisis (1969)     | The limitation that triggered AI Winter      |
| 01-08             | MLP Revival (1986)    | Backprop solves XOR + TinyDigits recognition |
| 01-09             | CNN Revolution (1998) | Convolutions for spatial understanding       |
| 01-08 + 11-13     | Transformers (2017)   | Structured sequence tasks with attention     |
| 01-08 + 14-19     | MLPerf (2018)         | Production optimization pipeline             |

See all milestones and their requirements:

```
tito milestone list
```

## 4.8 Quick Reference

Here are the commands you'll use throughout your journey:

```
# Module workflow
tito module start <N>          # Start working on module N
tito module complete <N>       # Export module to package
tito module status              # See your progress across all modules

# Milestones
tito milestone list             # See all milestones & requirements
tito milestone run <name>      # Run a milestone with your code

# Utilities
tito setup                     # First-time setup (safe to re-run)
tito system update              # Update TinyTorch (your work is preserved)
```



```
tito --help # Full command reference
```

## 4.9 Module Progression

TinyTorch has 20 modules organized in progressive tiers:

Table 4.2 groups the modules into tiers with scope and time estimates.

Table 4.2: TinyTorch module tiers with scope and time estimate.

| Tier                | Modules | Focus                                                            | Time Estimate |
|---------------------|---------|------------------------------------------------------------------|---------------|
| <b>Foundation</b>   | 01-08   | Core ML infrastructure (tensors, dataloader, autograd, training) | ~18-24 hours  |
| <b>Architecture</b> | 09-13   | Neural architectures (CNNs, transformers)                        | ~15-20 hours  |
| <b>Optimization</b> | 14-19   | Production optimization (profiling, quantization)                | ~18-24 hours  |
| <b>Capstone</b>     | 20      | Torch Olympics Competition                                       | ~8-10 hours   |

**Total: ~60-80 hours** over 14-18 weeks (4-6 hours/week pace).

See the module descriptions in this guide for detailed prerequisites and learning objectives.

### 4.10 Join the Community (Optional)

After setup, join the global TinyTorch community:

```
tito community login # Join the community
```

The community features include progress tracking and connecting with other builders.

### 4.11 For Instructors & TAs

#### **i** Note

TinyTorch's main path is self-paced learning. Instructors can also use the internal `tito nbgrader` workflow to stage notebooks for NBGrader-based courses.

**Instructor Workflow:**

- Stage TinyTorch notebooks as nbgrader source assignments
- Create release notebooks with solutions removed
- Collect and auto-grade submissions through nbgrader
- Export grades to CSV for LMS integration

Tito handles TinyTorch-specific staging; nbgrader handles release, collection, autograding, feedback, and grade export. Only content inside `### BEGIN SOLUTION / ### END SOLUTION` markers is treated as removable solution content during staging.

**Interested in early adoption?** [Join the discussion](#) to share your use case.

**Ready to start?** Run `tito module start 01` and begin building!

---

**PART I**

# **Foundation Tier**

---



## 🔥 Chapter 5

# Module 01: Tensor

Every byte of working memory your framework allocates flows through Tensor. Layout, strides, and dtype decide whether matmuls saturate the ALUs or stall on cache misses. Get the data structure right and the rest of the system scales. Get it wrong and no kernel tuning will save you.

### Module Info

**FOUNDATION TIER** | Difficulty: ●○○○ | Time: 4-6 hours | Prerequisites: None

**Prerequisites: None** means exactly that. This module assumes:

- Basic Python (lists, classes, methods)
- Basic math (matrix multiplication from linear algebra)
- No machine learning background required

If you can multiply two matrices by hand and write a Python class, you're ready.

## 5.1 Overview

Every neural network you have ever used — image classifiers, language models, self-driving perception stacks — is, at runtime, a sequence of operations on one data structure: the tensor. Get the tensor right and the rest of the framework practically writes itself. Get it wrong and every layer above leaks confusion.

In this module you build that data structure from scratch. By the end, your `Tensor` supports arithmetic, broadcasting, matrix multiplication, and shape manipulation — the same surface area you would call on `torch.Tensor`, backed by NumPy instead of CUDA.

## 5.2 Learning Objectives

### By completing this module, you will:

- **Implement** a complete Tensor class with arithmetic, matrix multiplication, shape manipulation, and reductions
- **Master** broadcasting semantics that enable efficient computation without data copying
- **Understand** computational complexity ( $O(n^3)$  for matmul) and memory trade-offs (views vs copies)
- **Connect** your implementation to production PyTorch patterns and design decisions

## 5.3 What You'll Build

### Implementation roadmap:

Table 5.1 lays out the implementation in order, one part at a time.




Figure 5.1: **Tensor Class Architecture:** One class packages data, shape metadata, and every operation the rest of TinyTorch will lean on.

Table 5.1: **Five-part implementation roadmap for the Tensor class.**

| Part | What You'll Implement                                                                         | Key Concept                                 |
|------|-----------------------------------------------------------------------------------------------|---------------------------------------------|
| 1    | <code>__init__</code> , <code>shape</code> , <code>size</code> , <code>dtype</code>           | Tensor as NumPy wrapper                     |
| 2    | <code>__add__</code> , <code>__sub__</code> , <code>__mul__</code> , <code>__truediv__</code> | Operator overloading + broadcasting         |
| 3    | <code>matmul()</code>                                                                         | Matrix multiplication with shape validation |
| 4    | <code>reshape()</code> , <code>transpose()</code>                                             | Shape manipulation, views vs copies         |
| 5    | <code>sum()</code> , <code>mean()</code> , <code>max()</code>                                 | Reductions along axes                       |

#### The pattern you'll enable:

```
# Computing predictions from data
output = x.matmul(W) + b # Matrix multiplication + bias (used in every neural
                           network)
```

### 5.3.1 What You're NOT Building

To keep this module focused, you will **not** implement:

- GPU support (NumPy runs on CPU only)
- Automatic differentiation
- Hundreds of tensor operations (PyTorch has 2000+, you'll build ~15 core ones)
- Memory optimization tricks (PyTorch uses lazy evaluation, memory pools, etc.)

**You are building the conceptual foundation.**

## 5.4 API Reference

This section provides a quick reference for the Tensor class you'll build. Think of it as your cheat sheet while implementing and debugging. Each method is documented with its signature and expected behavior.

### 5.4.1 Constructor

```
Tensor(data)
```

- data: list, numpy array, or scalar

### 5.4.2 Properties

Your Tensor wraps a NumPy array and exposes several properties that describe its structure. These properties are read-only and computed from the underlying data.

The properties you expose are summarised in Table 5.2.

Table 5.2: Read-only properties exposed by the Tensor class.

| Property | Type       | Description              |
|----------|------------|--------------------------|
| data     | np.ndarray | Underlying NumPy array   |
| shape    | tuple      | Dimensions, e.g., (2, 3) |
| size     | int        | Total number of elements |
| dtype    | np.dtype   | Data type (float32)      |

### 5.4.3 Arithmetic Operations

Python lets you override operators like + and \* by implementing special methods. When you write `x + y`, Python calls `x.__add__(y)`. Your implementations should handle both Tensor-Tensor operations and Tensor-scalar operations, letting NumPy's broadcasting do the heavy lifting.

The dunder methods you override map one-to-one onto Python operators, as Table 5.3 shows.

Table 5.3: Arithmetic dunder methods and their Python operator equivalents.

| Operation      | Method                   | Example                                  |
|----------------|--------------------------|------------------------------------------|
| Addition       | <code>__add__</code>     | <code>x + y</code> or <code>x + 2</code> |
| Subtraction    | <code>__sub__</code>     | <code>x - y</code>                       |
| Multiplication | <code>__mul__</code>     | <code>x * y</code>                       |
| Division       | <code>__truediv__</code> | <code>x / y</code>                       |

### 5.4.4 Matrix & Shape Operations

These methods transform tensors without changing their data (for views) or perform mathematical operations that produce new data (for matmul).

Table 5.4 lists the shape-transforming methods.

Table 5.4: Matrix and shape manipulation methods on Tensor.

| Method                 | Signature                                                 | Description                            |
|------------------------|-----------------------------------------------------------|----------------------------------------|
| <code>matmul</code>    | <code>matmul(other) -&gt; Tensor</code>                   | Matrix multiplication                  |
| <code>reshape</code>   | <code>reshape(*shape) -&gt; Tensor</code>                 | Change shape (-1 to infer)             |
| <code>transpose</code> | <code>transpose(dim0=None, dim1=None) -&gt; Tensor</code> | Swap dimensions (defaults to last two) |

### 5.4.5 Reductions

Reduction operations collapse one or more dimensions by aggregating values. The `axis` parameter controls which dimension gets collapsed. If `axis=None`, all dimensions collapse to a single scalar.

Table 5.5 lists the reduction methods and their axis semantics.

Table 5.5: Reduction methods that collapse tensor dimensions.

| Method            | Signature                                                 | Description      |
|-------------------|-----------------------------------------------------------|------------------|
| <code>sum</code>  | <code>sum(axis=None, keepdims=False) -&gt; Tensor</code>  | Sum elements     |
| <code>mean</code> | <code>mean(axis=None, keepdims=False) -&gt; Tensor</code> | Average elements |
| <code>max</code>  | <code>max(axis=None, keepdims=False) -&gt; Tensor</code>  | Maximum element  |

## 5.5 Core Concepts

Five ideas drive every line of code in this module: dimensionality, broadcasting, views vs. copies, matrix multiplication, and reduction along an axis. They are not TinyTorch-specific. Internalize them here and you will recognize them in PyTorch, JAX, NumPy, and every framework you touch afterwards.

### 5.5.1 Tensor Dimensionality

Tensors generalize the shapes you already know. A scalar is a single number — a temperature reading of 72.5 degrees. Stack scalars into a list and you get a vector — a day’s worth of readings. Stack vectors and you get a matrix — a spreadsheet of readings across many days. Keep stacking and you reach 3D and 4D tensors that hold video frames, batches of images, or sequences of token embeddings.

One class handles all of these cases. The same code path that adds two scalars adds two 4D tensors, because broadcasting (next section) makes shape compatibility a runtime concern, not a compile-time one.

Table 5.6 anchors each rank in a concrete example you have seen before.

Table 5.6: Tensor ranks from scalar to 4D with concrete examples.

| Rank | Name   | Shape                   | Concrete Example                           |
|------|--------|-------------------------|--------------------------------------------|
| 0D   | Scalar | <code>()</code>         | Temperature reading: 72.5                  |
| 1D   | Vector | <code>(768,)</code>     | Audio sample: 768 measurements             |
| 2D   | Matrix | <code>(128, 768)</code> | Spreadsheet: 128 rows $\times$ 768 columns |



| Rank | Name      | Shape             | Concrete Example                    |
|------|-----------|-------------------|-------------------------------------|
| 3D   | 3D Tensor | (32, 224, 224)    | Video frames: 32 grayscale images   |
| 4D   | 4D Tensor | (32, 3, 224, 224) | Video frames: 32 color (RGB) images |

### 5.5.2 Broadcasting

When you add a vector to a matrix, the shapes don't match. Should this fail? In most programming contexts, yes. But many computations need to apply the same operation across rows or columns. For example, if you want to adjust all values in a spreadsheet by adding a different offset to each column, you need to add a vector to a matrix. NumPy and PyTorch implement broadcasting to handle this: automatically expanding smaller tensors to match larger ones without actually copying data.

Consider adding a bias vector `[10, 20, 30]` to every row of a matrix. Without broadcasting, you'd need to manually tile the vector into a matrix first, wasting memory. With broadcasting, the operation just works, and the framework handles alignment internally.

Here is what `__add__` looks like once broadcasting does the work for you:

```
def __add__(self, other):
    """Add two tensors element-wise with broadcasting support."""
    if isinstance(other, Tensor):
        return Tensor(self.data + other.data) # NumPy handles broadcasting
    else:
        return Tensor(self.data + other)      # Scalar broadcast
```

There is no shape logic in this method. NumPy aligns the operand shapes from right to left, expands dimensions where one side is 1, and refuses the operation otherwise. You inherit that contract for free.




Figure 5.2: **Broadcasting Mechanics:** How a vector is expanded to match a matrix's shape during arithmetic operations.

The rules are simpler than they look. Compare shapes from right to left. At each position, dimensions are compatible if they're equal or if one of them is 1. Missing dimensions on the left are treated as 1. If any position fails this check, broadcasting fails.

Table 5.7 walks through the compatibility check on representative shape pairs.

Table 5.7: **Broadcasting compatibility for representative shape pairs.**

| Shape A   | Shape B | Result    | Valid?       |
|-----------|---------|-----------|--------------|
| (3, 4)    | (4, )   | (3, 4)    |              |
| (3, 4)    | (3, 1)  | (3, 4)    |              |
| (3, 4)    | (3, )   | Error     | (3 $\neq$ 4) |
| (2, 3, 4) | (3, 4)  | (2, 3, 4) |              |

The memory savings are not academic. Adding a (768, ) vector to a (32, 512, 768) tensor would, without broadcasting, require tiling the vector 32×512 times — **48 MB** of redundant data, about **12.6 million** float32 numbers. With broadcasting, you store the original **3 KB** vector and read it through expanded indices.

### 5.5.3 Views vs. Copies

When you reshape a tensor, does it allocate new memory or just create a different view of the same data? The answer has huge implications for both performance and correctness.

A view shares memory with its source. Reshaping a 1 GB tensor is instant because you’re just changing the metadata that describes how to interpret the bytes, not copying the bytes themselves. But this creates an important gotcha: modifying a view modifies the original.

```
x = Tensor([1, 2, 3, 4])
y = x.reshape(2, 2) # y is a VIEW of x
y.data[0, 0] = 99   # This also changes x!
```

Arithmetic operations like addition always create copies because they compute new values. This is safer but uses more memory. Production code carefully manages views to avoid both memory blowup (too many copies) and silent bugs (unexpected mutations through views).

Table 5.8 contrasts the memory and time cost of view operations against the copying ones.

Table 5.8: **Memory and time cost of view operations versus copying ops.**

| Operation                                                   | Type  | Memory         | Time |
|-------------------------------------------------------------|-------|----------------|------|
| <code>reshape()</code>                                      | View* | Shared         | O(1) |
| <code>transpose()</code>                                    | View* | Shared         | O(1) |
| <code>+</code> <code>-</code> <code>*</code> <code>/</code> | Copy  | New allocation | O(n) |

\*When data is contiguous in memory

### 5.5.4 Matrix Multiplication

Matrix multiplication is the computational workhorse of neural networks. Every linear layer, every attention head, every embedding lookup involves matmul. Understanding its mechanics and cost is essential.

The operation is simple in concept: for each output element, compute a dot product of a row from the first matrix with a column from the second. But this simplicity hides cubic complexity. Multiplying two  $n \times n$  matrices requires  $n^3$  multiplications and  $n^3$  additions.

Here’s how the educational implementation in your module works:

The code in `#lst-01-tensor-matmul` makes this concrete.

```
def matmul(self, other):
    """Matrix multiplication of two tensors."""
    if not isinstance(other, Tensor):
        raise TypeError(f"Expected Tensor for matrix multiplication, got {type(other)}")

    # Shape validation: inner dimensions must match
    if len(self.shape) >= 2 and len(other.shape) >= 2:
        if self.shape[-1] != other.shape[-2]:
            raise ValueError(
                f"Cannot perform matrix multiplication: {self.shape} @ {other.shape}. "
                f"Inner dimensions must match: {self.shape[-1]} ≠ {other.shape[-2]}"
            )

    a = self.data
    b = other.data

    # Handle 2D matrices with explicit loops (educational)
    if len(a.shape) == 2 and len(b.shape) == 2:
        M, K = a.shape
        K2, N = b.shape
        result_data = np.zeros((M, N), dtype=a.dtype)

        # Each output element is a dot product
        for i in range(M):
            for j in range(N):
                result_data[i, j] = np.dot(a[i, :], b[:, j])
    else:
        # For batched operations, use np.matmul
        result_data = np.matmul(a, b)

    return Tensor(result_data)
```

: Listing 1.1 — Matrix multiplication with explicit loops for the 2D case and shape validation. `{#lst-01-tensor-matmul}`

The explicit loops in the 2D case are intentionally slower than `np.matmul` because they reveal exactly what matrix multiplication does: each output element requires  $K$  operations, and there are  $M \times N$  outputs, giving  $O(M \times K \times N)$  total operations. For square matrices, this is  $O(n^3)$ .

### **i** Systems Implication: Cache Tiling

While element-wise operations like addition are **Memory-Bound** (hardware sits idle waiting for data from RAM), Matrix Multiplication is heavily **Compute-Bound** ( $O(n^3)$  math for  $O(n^2)$  bytes). A naive nested `for` loop constantly evicts data from the ultra-fast L1 cache back to slow main memory. Optimized libraries like `np.matmul` (which calls BLAS) use *Cache Tiling*—breaking matrices into tiny blocks that fit perfectly into the L1 cache, keeping the ALUs fed without waiting for RAM.

### 5.5.5 Shape Manipulation

Shape manipulation operations change how data is interpreted without changing the values themselves. Understanding when data is copied versus viewed is crucial for both correctness and performance.

The `reshape` method reinterprets the same data with different dimensions:

The code in `?lst-01-tensor-reshape` makes this concrete.

```
def reshape(self, *shape):
    """Reshape tensor to new dimensions."""
    if len(shape) == 1 and isinstance(shape[0], (tuple, list)):
        new_shape = tuple(shape[0])
    else:
        new_shape = shape

    # Handle -1 for automatic dimension inference
    if -1 in new_shape:
        known_size = 1
        unknown_idx = new_shape.index(-1)
        for i, dim in enumerate(new_shape):
            if i != unknown_idx:
                known_size *= dim
        unknown_dim = self.size // known_size
        new_shape = list(new_shape)
        new_shape[unknown_idx] = unknown_dim
        new_shape = tuple(new_shape)

    # Validate total elements match
    if np.prod(new_shape) != self.size:
        target_size = int(np.prod(new_shape))
        raise ValueError(
            f"Cannot reshape {self.shape} to {new_shape}\n"
            f"    Element count mismatch: {self.size} elements vs {target_size} elements\n"
            f"    Reshape preserves data, so total elements must stay the same\n"
            f"    Use -1 to infer a dimension: reshape(-1, {new_shape[-1]} if len(new_shape) > 0 else 1}) lets NumPy calculate"
        )

    reshaped_data = np.reshape(self.data, new_shape)
    return Tensor(reshaped_data, requires_grad=self.requires_grad)
```

: Listing 1.2 — Reshape with dimension inference via `-1` and element-count validation. `{#lst-01-tensor-reshape}`

The `-1` syntax is particularly useful: it tells NumPy to infer one dimension automatically. When flattening a batch of images, `x.reshape(batch_size, -1)` lets NumPy calculate the feature dimension.

While these mathematical abstractions are powerful, they mask the physical reality of how data is stored and accessed. Understanding the gap between mathematical shape and physical memory layout is what separates a machine learning practitioner from a systems engineer.

**i Systems Implication: Contiguous Memory & Strides**

To the math, `reshape` or `transpose` just change dimensions. To the hardware, they manipulate *strides*. A tensor's data is stored as a flat, 1D contiguous block of bytes in RAM. A `reshape` is an  $O(1)$  metadata update—the underlying memory is completely untouched.

However, `transpose` introduces a systems trap. By swapping dimensions, you swap strides. The memory is no longer sequentially aligned with the new shape. If you iterate over a transposed tensor, the processor jumps back and forth across RAM, destroying spatial locality and causing massive **cache misses**. This is why systems engineers call `.contiguous()` after a `transpose`—trading a one-time  $O(n)$  memory copy to restore sequential alignment and keep the cache fed.

### 5.5.6 Computational Complexity

As we saw with the explicit loops in our matrix multiplication implementation, not all tensor operations are equal. Element-wise operations like addition visit each element exactly once, resulting in  $O(n)$  time complexity where  $n$  is the total number of elements. Reductions like `sum` also visit each element once. But matrix multiplication is fundamentally different.

Multiplying two  $n \times n$  matrices requires  $n^3$  operations: for each of the  $n^2$  output elements, you compute a dot product of  $n$  values. This cubic scaling is why a  $2000 \times 2000$  matmul takes  $8 \times$  longer than a  $1000 \times 1000$  matmul, not  $4 \times$ . In neural networks, matrix multiplications consume over 90% of compute time. This computational density is precisely why GPUs exist for ML: a modern GPU has thousands of cores that can compute thousands of dot products simultaneously, turning an 800ms CPU operation into an 8ms GPU operation.

Table 5.9 summarises the time and memory cost of the core operations.

Table 5.9: Computational complexity of core tensor operations.

| Operation                | Complexity | Notes                    |
|--------------------------|------------|--------------------------|
| Element-wise (+, -, *)   | $O(n)$     | Linear in tensor size    |
| Reductions (sum, mean)   | $O(n)$     | Must visit every element |
| Matrix multiply (matmul) | $O(n^3)$   | Dominates training time  |

### 5.5.7 Axis Semantics

The `axis` parameter in reductions specifies which dimension to collapse. Think of it as “sum along this axis” or “average out this dimension.” The result has one fewer dimension than the input.

For a 2D tensor with shape (rows, columns), summing along axis 0 collapses the rows, giving you column totals. Summing along axis 1 collapses the columns, giving you row totals. Summing with `axis=None` collapses everything to a single scalar.

Your reduction implementations simply pass the axis to NumPy:

```
def sum(self, axis=None, keepdims=False):
    """Sum tensor along specified axis."""
    result = np.sum(self.data, axis=axis, keepdims=keepdims)
    return Tensor(result)
```

The `keepdims=True` option preserves the reduced dimension as size 1, which is useful for broadcasting the result back.

For shape (rows, columns) = (32, 768):

```
sum(axis=0) → collapse rows    → shape (768,) - column totals
sum(axis=1) → collapse columns → shape (32,)  - row totals
sum(axis=None) → collapse all  → scalar
```

Visual:

```
[[1, 2, 3],      sum(axis=0)    sum(axis=1)
 [4, 5, 6]] →    [5, 7, 9]    or [6, 15]
                (down cols)    (across rows)
```

## 5.6 Architecture

Your `Tensor` sits at the top of a stack that reaches down to hardware. When you call `x + y`, Python calls your `__add__` method, which delegates to NumPy, which calls optimized BLAS libraries written in C and Fortran, which use CPU SIMD instructions that process multiple numbers in a single clock cycle.




Figure 5.3: **The Tensor Stack:** From the Python interface down to the hardware CPU instructions.

This is the same architecture used by PyTorch and TensorFlow, just with different backends. PyTorch replaces NumPy with a C++ engine and BLAS with CUDA kernels running on GPUs. But the Python interface and the abstractions are identical. When you understand TinyTorch’s `Tensor`, you understand PyTorch’s `Tensor`.

### 5.6.1 Module Integration

Your `Tensor` is the only data structure the rest of TinyTorch carries around. Every subsequent module either operates on tensors, returns tensors, or stores tensors as state.

```
Module 01: Tensor      ← you are here
  → Module 02: Activations (element-wise functions on tensors)
  → Module 03: Layers      (parameters stored as tensors)
  → Module 06: Autograd    (gradients tracked alongside tensors)
  → Module 08: Training    (every step reads and writes tensors)
```

Bugs you ship from this module surface as confusing failures three modules later. The investment in clean shape handling, consistent dtypes, and correct broadcasting pays off through the entire stack.

## 5.7 Common Errors & Debugging

Five errors account for almost every confused half-hour you will spend in this module. Read them now so you recognize them when they appear, instead of debugging them blind.

### 5.7.1 Shape Mismatch in matmul

**Error:** `ValueError: shapes (2,3) and (2,2) not aligned`

Matrix multiplication requires the inner dimensions to match. If you're multiplying  $(M, K)$  by  $(K, N)$ , both  $K$  values must be equal. The error above happens when trying to multiply a  $(2,3)$  matrix by a  $(2,2)$  matrix:  $3 \neq 2$ .

**Fix:** Check your shapes. The rule is `a.shape[-1]` must equal `b.shape[-2]`.

### 5.7.2 Broadcasting Failures

**Error:** `ValueError: operands could not be broadcast together`

Broadcasting fails when shapes can't be aligned according to the rules. Remember: compare right to left, and dimensions must be equal or one must be 1.

**Examples:** -  $(2, 3) + (3, )$  works - 3 matches 3, and the missing dimension becomes 1 -  $(2, 3) + (2, )$  fails - comparing right to left:  $3 \neq 2$

**Fix:** Reshape to make dimensions compatible: `vector.reshape(-1, 1)` or `vector.reshape(1, -1)`

### 5.7.3 Reshape Size Mismatch

**Error:** `ValueError: cannot reshape array of size X into shape Y`

Reshape only rearranges elements; it can't create or destroy them. If you have 12 elements, you can reshape to  $(3, 4)$  or  $(2, 6)$  or  $(2, 2, 3)$ , but not to  $(5, 5)$ .

**Fix:** Ensure `np.prod(old_shape) == np.prod(new_shape)`

### 5.7.4 Missing Attributes

**Error:** `AttributeError: 'Tensor' has no attribute 'shape'`

Your `__init__` method needs to set all required attributes. If you forget to set `self.shape`, any code that accesses `tensor.shape` will fail.

**Fix:** Add `self.shape = self.data.shape` in your constructor

### 5.7.5 Type Errors in Arithmetic

**Error:** `TypeError: unsupported operand type(s) for +: 'Tensor' and 'int'`

Your arithmetic methods need to handle both Tensor and scalar operands. When someone writes `x + 2`, your `__add__` receives the integer 2, not a Tensor.

**Fix:** Check for scalars: `if isinstance(other, (int, float)): ...`

## 5.8 Production Context

### 5.8.1 Your Implementation vs. PyTorch

Your `Tensor` and PyTorch's `torch.Tensor` share the same conceptual design. The differences are in the engine room: PyTorch swaps NumPy for a C++ runtime, runs on CUDA, Metal, or ROCm, and exposes

thousands of specialized kernels. The Python-facing API, the broadcasting rules, and the shape semantics are the same.

Table 5.10 places your implementation side by side with the production reference for direct comparison.

Table 5.10: Feature comparison between TinyTorch Tensor and PyTorch Tensor.

| Feature           | Your Implementation | PyTorch           |
|-------------------|---------------------|-------------------|
| <b>Backend</b>    | NumPy (Python)      | C++/CUDA          |
| <b>Speed</b>      | 1x (baseline)       | 10-100x faster    |
| <b>GPU</b>        | CPU only            | CUDA, Metal, ROCm |
| <b>Operations</b> | ~15 core ops        | 2000+ operations  |

### 5.8.2 Code Comparison

The following comparison shows equivalent operations in TinyTorch and PyTorch. Notice how closely the APIs mirror each other. This is intentional: by learning TinyTorch's patterns, you're simultaneously learning PyTorch's patterns.

## 5.9 Your TinyTorch

```
from tinytorch.core.tensor import Tensor

x = Tensor([[1, 2], [3, 4]])
y = x + 2
z = x.matmul(w)
loss = z.mean()
```

## 5.10 PyTorch

```
import torch

x = torch.tensor([[1, 2], [3, 4]], dtype=torch.float32)
y = x + 2
z = x @ w
loss = z.mean()
```

Let's walk through each line to understand the comparison:

- **Line 1 (Import):** Both frameworks use a simple import. TinyTorch exposes `Tensor` from `core.tensor`; PyTorch uses `torch.tensor()` as a factory function.
- **Line 3 (Creation):** TinyTorch infers `dtype` from input; PyTorch requires explicit `dtype=torch.float32` for floating-point operations. This explicitness matters for performance tuning in production.
- **Line 4 (Broadcasting):** Both handle `x + 2` identically, broadcasting the scalar across all elements. Same semantics, same result.
- **Line 5 (Matrix multiply):** TinyTorch uses `.matmul()` method; PyTorch supports both `.matmul()` and the `@` operator. The operation is identical.



- **Line 6 (Reduction):** Both use `.mean()` to reduce the tensor to a scalar. Reductions like this are fundamental to computing loss values.

#### 💡 What's Identical

Broadcasting rules, shape semantics, and API design patterns. When you debug PyTorch shape errors, you'll understand exactly what's happening because you built the same abstractions.

### 5.10.1 Why Tensors Matter at Scale

The class you are building looks small. The systems built on it do not. Three numbers set the scale:

- **Large language models** — A 175B-parameter model stored in fp16 is **326 GB** of tensor data. That is one model, one copy, before optimizer state.
- **Image batches** — A batch of 128 RGB images at  $224 \times 224$  in float32 is **73.5 MB**. A training run streams thousands of these per minute.
- **Self-driving perception** — Multi-camera pipelines run tensor ops at **36 FPS**, so each frame's millions of operations have to finish in 28 ms.

A single matrix multiply can absorb 90% of training time. Every micro-optimization in a real framework — contiguous strides, fused kernels, view-vs-copy decisions — exists because the `Tensor` you are about to write multiplies its mistakes by billions.

## 5.11 Check Your Understanding

### 💡 Check Your Understanding — Tensor

Before moving on, verify you can articulate each of the following:

- Why `reshape` is  $O(1)$  metadata update while `transpose` breaks contiguity and triggers cache misses on later iteration.
- How right-to-left broadcasting lets  $(32, 1, 768) + (512, 768)$  produce  $(32, 512, 768)$  without materializing a tiled copy.
- Why `matmul` is  $O(n^3)$  — and why doubling matrix size multiplies runtime by 8x, not 4x.
- The memory distinction between views (share bytes with source) and copies (fresh allocation), and why silent mutation through a view is a systems bug.
- How a scalar logit of  $32 \times 3 \times 224 \times 224$  image batch lands at ~18.4 MB in float32, and why batch size is the first knob to turn when the GPU runs Out of Memory (OOM).

If any of these feels fuzzy, revisit Core Concepts (Broadcasting, Views vs. Copies, Matrix Multiplication, Computational Complexity) before moving on.

The collapsible Q&A below works through each of these in depth — skim the ones that still feel fuzzy.

#### Q1: Memory Calculation

A batch of 32 RGB images ( $224 \times 224$  pixels) stored as float32. How much memory?

#### 💡 Answer

$32 \times 3 \times 224 \times 224 \times 4 = 19,267,584 \text{ bytes} \approx 18.4 \text{ MB}$

Double the batch, double the memory — which is why batch size is the first knob people turn when training runs out of GPU memory.

#### Q2: Broadcasting Savings

Adding a vector  $(768,)$  to a 3D tensor  $(32, 512, 768)$ . How much memory does broadcasting save?

#### 💡 Answer

Without broadcasting:  $32 \times 512 \times 768 \times 4 = 48.0 \text{ MB}$

With broadcasting:  $768 \times 4 = 3 \text{ KB}$

Savings: **~48 MB per operation**. A transformer layer performs hundreds of broadcasts per forward pass; the cost of getting this wrong scales with the model.

#### Q3: Matmul Scaling

If a  $1000 \times 1000$  matmul takes 100ms, how long will  $2000 \times 2000$  take?

#### 💡 Answer

Matmul is  $O(n^3)$ . Doubling  $n \rightarrow 2^3 = 8x \text{ longer} \rightarrow \sim 800\text{ms}$

This is why matrix size matters so much for transformer scaling!

#### Q4: Shape Prediction

What's the output shape of  $(32, 1, 768) + (512, 768)$ ?

#### 💡 Answer

Broadcasting aligns right-to-left:  $(32, 1, 768) - (512, 768)$

Result:  **$(32, 512, 768)$**

The 1 broadcasts to 512, and 32 is prepended.

#### Q5: Views vs Copies

You reshape a 1GB tensor, then modify one element in the reshaped version. What happens to the original tensor? What if you had used  $x + 0$  instead of reshape?

#### 💡 Answer

**Reshape (view):** The original tensor IS modified. Reshape creates a view that shares memory with the original. Changing `y.data[0,0] = 99` also changes `x.data[0]`.

**Addition (copy):** The original tensor is NOT modified.  $x + 0$  creates a new tensor with freshly allocated memory. The values are identical but stored in different locations.

This distinction matters enormously for:

- **Memory** — Views use 0 extra bytes; copies use  $n$  extra bytes.
- **Performance** — Views are  $O(1)$ ; copies are  $O(n)$ .
- **Correctness** — Unexpected mutations through views are a common source of bugs.

## 5.12 Key Takeaways

- **Tensor as universal carrier:** one class holds every value the rest of the framework touches — scalars, vectors, 4-D image batches — because broadcasting makes shape compatibility a runtime concern, not a compile-time one.
- **Broadcasting replaces tiling:** right-to-left alignment lets a  $(768,)$  bias merge with a  $(32, 512, 768)$  activation tensor for the cost of the vector, not the expanded 48 MB copy.

- **Views vs. copies is a systems decision:** `reshape` and `transpose` are  $O(1)$  metadata edits that share memory; arithmetic ops allocate new buffers. Mixing them up causes either memory blowup or silent mutation.
- **Matmul is the cubic tax:**  $O(n^3)$  compute on  $O(n^2)$  bytes makes matmul compute-bound and cache-sensitive, which is why every real framework routes it through tiled BLAS kernels.
- **Contiguity is physical, shape is logical:** `transpose` rearranges strides without moving bytes, so iteration order breaks cache locality — hence the `.contiguous()` call after transpose in production code.

**Coming next:** Module 02 adds nonlinearity — ReLU, Sigmoid, Tanh, GELU, Softmax — as element-wise functions that consume and return the `Tensor` you just built.

## 5.13 Further Reading

The abstractions we rely on today—broadcasting, memory views, and hardware-accelerated matrix multiplication—were not developed overnight. They are the culmination of decades of research bridging mathematics and computer architecture. For students who want to trace the evolution of these ideas from early linear algebra libraries to modern tensor frameworks, the following literature is essential.

### 5.13.1 Seminal Papers

- **NumPy: Array Programming** - Harris et al. (2020). The definitive reference for NumPy, which underlies your `Tensor` implementation. Explains broadcasting, views, and the design philosophy. **Systems Implication:** Standardized memory layouts (strides) and contiguous blocks allowed C-level operations to bypass the slow Python interpreter, maximizing memory bandwidth. [Nature](#)
- **BLAS (Basic Linear Algebra Subprograms)** - Lawson et al. (1979). The foundation of all high-performance matrix operations. Your `np.matmul` ultimately calls BLAS routines optimized over 40+ years. Understanding BLAS levels (1, 2, 3) explains why matmul is special. **Systems Implication:** Defined the memory hierarchy abstractions (registers, L1/L2 cache, RAM) allowing matrix multiplication to be tiled for optimal cache reuse, fundamentally defining modern dense compute limits. [ACM TOMS](#)
- **Automatic Differentiation in ML** - Baydin et al. (2018). Survey of automatic differentiation techniques. **Systems Implication:** Reverse-mode AD required materializing the intermediate forward pass tensors in memory, making memory capacity a primary bottleneck for deep learning before checkpointing techniques were developed. [JMLR](#)

### 5.13.2 Additional Resources

- **Textbook:** [“Deep Learning”](#) by Goodfellow, Bengio, and Courville - Chapter 2 covers linear algebra foundations including tensor operations
- **Documentation:** [PyTorch Tensor Tutorial](#) - See how production frameworks implement similar concepts

## 5.14 What’s Next?

You now have the universal carrier of every value the rest of the framework will ever compute. A `Tensor` holds the data, knows its shape, and supports the arithmetic that linear algebra and eventually backpropagation depend on.

The next module is **Module 02 — Activations**. It answers a single question: *given a tensor of pre-activations, how do you apply a non-linearity element-wise — and why is that the one operation that lets a deep network represent anything beyond a glorified linear regression?* You will implement ReLU, Sigmoid, Tanh, and Softmax on top of the `Tensor` you just built, treating each one as a pure function from tensor to tensor. Everything works

because `__add__`, `__mul__`, broadcasting, and shape preservation are already in place. That is the dividend of getting Module 01 right.

## 5.15 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.

## 🔥 Chapter 6

# Module 02: Activations

Activations are the first operator in your framework with zero arithmetic intensity: every element is touched once and reused zero times. That makes them pure memory-bandwidth workloads. Stacking them also breaks the linear collapse that would otherwise flatten a hundred layers into one matrix multiply, so both the math and the memory-wall lesson start here.

**i** Module Info

**FOUNDATION TIER** | Difficulty: ●○○○ | Time: 3-5 hours | Prerequisites: 01 (Tensor)

**Prerequisites: Module 01 (Tensor)** means you need:

- Completed Tensor implementation with element-wise operations
- Understanding of tensor shapes and broadcasting
- Familiarity with NumPy mathematical functions

If you can create a Tensor and perform element-wise arithmetic ( $x + y$ ,  $x * 2$ ), you're ready.

## 6.1 Overview

A neural network without activation functions isn't a neural network — it's a single matrix multiplication wearing a costume. Stack a hundred linear layers, and the composition collapses to one:  $W_2(W_1x) = (W_2W_1)x$ . Depth buys you nothing until you break the linearity.

Activations are how you break it. ReLU zeros out negatives. Sigmoid squashes any real number into  $(0, 1)$ . Softmax turns raw scores into a probability distribution. Each is just a few lines of math, but together they're what lets a network learn to distinguish a cat from a dog instead of computing one giant linear regression.

You'll build five of them: ReLU, Sigmoid, Tanh, GELU, and Softmax. Along the way you'll meet the chapter's load-bearing insight — *why every production softmax subtracts the max before exponentiating* — and the dead-neuron problem that explains why ReLU's apparent simplicity hides a real failure mode.

## 6.2 Learning Objectives

**💡** By completing this module, you will:

- **Implement** five core activation functions (ReLU, Sigmoid, Tanh, GELU, Softmax) with the numerical-stability tricks production frameworks use
- **Explain** why nonlinearity turns a stack of matrix multiplies into a function approximator
- **Quantify** the compute cost of each activation and decide when the extra accuracy of GELU is worth the extra exponentials
- **Map** your implementations onto the corresponding `torch.nn.functional` calls so PyTorch stops feeling like a black box

## 6.3 What You'll Build




Figure 6.1: **Activation functions in TinyTorch:** ReLU, Sigmoid, Tanh, GELU, and Softmax transformations.

### Implementation roadmap:

Table 6.1 lays out the implementation in order, one part at a time.

Table 6.1: **Implementation roadmap for the core activation functions.**

| Part | What You'll Implement          | Key Concept                                        |
|------|--------------------------------|----------------------------------------------------|
| 1    | <code>ReLU.forward()</code>    | Sparsity through zeroing negatives                 |
| 2    | <code>Sigmoid.forward()</code> | Mapping to (0,1) for probabilities                 |
| 3    | <code>Tanh.forward()</code>    | Zero-centered activation for better gradients      |
| 4    | <code>GELU.forward()</code>    | Smooth nonlinearity for transformers               |
| 5    | <code>Softmax.forward()</code> | Probability distributions with numerical stability |

### The pattern you'll enable:

```

# Transforming tensors through nonlinear functions
relu = ReLU()
activated = relu(x)  # Zeros negatives, keeps positives

softmax = Softmax()
probabilities = softmax(logits)  # Converts to probability distribution (sums to 1)

```

### 6.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Gradient computation (`backward()` methods are stubs for now — automatic differentiation is a later module)
- Learnable parameters (activations are fixed mathematical functions)
- Advanced variants (LeakyReLU, ELU, Swish — PyTorch ships dozens; you'll build the core five that cover ~95% of real architectures)
- GPU acceleration (your NumPy implementation runs on CPU)

The forward pass is enough to build intuition. Gradients show up in Module 06.

## 6.4 API Reference

This section provides a quick reference for the activation classes you'll build. Each activation is a callable object with a `forward()` method that transforms an input tensor element-wise.

### 6.4.1 Activation Pattern

All activations follow this structure:

```
class ActivationName:
    def forward(self, x: Tensor) -> Tensor:
        # Apply mathematical transformation
        pass

    def __call__(self, x: Tensor) -> Tensor:
        return self.forward(x)

    def backward(self, grad: Tensor) -> Tensor:
        # Stub - autograd adds gradient computation later
        pass
```

### 6.4.2 Core Activations

Table 6.2 gives the mathematical form, output range, and typical use case of each.

Table 6.2: Mathematical form, output range, and use case for each activation.

| Activation | Mathematical Form                 | Output Range           | Primary Use Case             |
|------------|-----------------------------------|------------------------|------------------------------|
| ReLU       | $\max(0, x)$                      | $[0, \infty)$          | Hidden layers (CNNs, MLPs)   |
| Sigmoid    | $1 / (1 + e^{-x})$                | $(0, 1)$               | Binary classification output |
| Tanh       | $(e^x - e^{-x}) / (e^x + e^{-x})$ | $(-1, 1)$              | RNNs, zero-centered needs    |
| GELU       | $x \cdot \Phi(x)$                 | $(-\infty, \infty)$    | Transformers (GPT, BERT)     |
| Softmax    | $e^{x_i} / \sum e^{x_j}$          | $(0, 1), \text{sum}=1$ | Multi-class classification   |

### 6.4.3 Method Signatures

#### ReLU

```
ReLU.forward(x: Tensor) -> Tensor
```

Sets negative values to zero, preserves positive values.

#### Sigmoid

```
Sigmoid.forward(x: Tensor) -> Tensor
```

Maps any real number to (0, 1) range using logistic function.

#### Tanh

```
Tanh.forward(x: Tensor) -> Tensor
```

Maps any real number to (-1, 1) range using hyperbolic tangent.

#### GELU

```
GELU.forward(x: Tensor) -> Tensor
```

Smooth approximation to ReLU using Gaussian error function.

#### Softmax

```
Softmax.forward(x: Tensor, dim: int = -1) -> Tensor
```

Converts vector to probability distribution along specified dimension.

## 6.5 Core Concepts

Before you write a single `np.maximum`, settle the four ideas that govern every activation choice you'll make for the rest of the book: why nonlinearity is non-optional, why ReLU dominates in practice, why softmax has to be implemented carefully, and what each activation costs you per element.

### 6.5.1 Why Non-linearity Matters

Consider what happens when you stack linear transformations. If you multiply a matrix by a vector, then multiply the result by another matrix, the composition is still just matrix multiplication. Mathematically:

$$f(x) = W(Wx) = (WW)x = Wx$$

A 100-layer network of pure matrix multiplications is identical to a single matrix multiplication. The depth buys you nothing.

Activation functions break this linearity. When you insert  $f(x) = \max(0, x)$  between layers, the composition becomes nonlinear:

$$f(x) = \max(0, W \max(0, Wx))$$

Now you can't simplify the layers away. Each layer learns to detect a different level of structure. Layer 1 finds edges; layer 2 composes edges into shapes; layer 3 composes shapes into objects. This hierarchy of features is *only* possible because the activation between layers is nonlinear — the moment you remove it, the whole stack collapses back into a single matrix and the hierarchy vanishes with it.

That's the entire reason this module exists. Five small functions are what separate a neural network from a linear regression with extra steps.



### 6.5.2 ReLU and Its Variants

ReLU (Rectified Linear Unit) is deceptively simple: it zeros out negative values and leaves positive values unchanged. Here's the complete implementation from your module:

```
class ReLU:
    def forward(self, x: Tensor) -> Tensor:
        """Apply ReLU activation element-wise."""
        result = np.maximum(0, x.data)
        return Tensor(result)
```

This simplicity is ReLU's greatest strength. The operation is a single comparison per element:  $O(n)$  with a tiny constant factor. Modern CPUs can execute billions of comparisons per second. Compare this to sigmoid, which requires computing an exponential for every element.

ReLU creates **sparsity**. When half of your activations are exactly zero, computations become faster (multiplying by zero is free) and models generalize better (sparse representations are less prone to overfitting). In a 1000-neuron layer, ReLU typically activates 300-500 neurons, effectively creating a smaller, specialized network for each input.

The discontinuity at zero is both a feature and a bug. ReLU's gradient is exactly 1 for positive inputs and exactly 0 for negative inputs — no decay, no saturation, no vanishing. That single property is what made training 100-layer networks tractable.

The bug is a corollary: **the dying ReLU problem**. The failure chain is short and brutal:

1. A weight update pushes a neuron's pre-activation negative for *every* input in the batch.
2. ReLU outputs zero.
3. The local gradient is zero.
4. Backprop sends zero through it, so its weights never update again.

The neuron is now a permanent zero. It contributes nothing to the forward pass and receives nothing from the backward pass. In practice, 10–40% of ReLU units in a trained network are dead. This is what variants like LeakyReLU, ELU, and GELU were designed to fix — they keep a small nonzero gradient on the negative side so a struggling neuron can claw its way back. Despite this, plain ReLU remains the default in CNNs and MLPs: it's fast, it doesn't vanish, and a few dead neurons in an over-parameterized model are cheaper than the extra exponentials.

### 6.5.3 Sigmoid and Tanh

Sigmoid maps any real number to the range  $(0, 1)$ , making it perfect for representing probabilities:

The code in `?@lst-02-activations-sigmoid` makes this concrete.

```
class Sigmoid:
    def forward(self, x: Tensor) -> Tensor:
        """Apply sigmoid activation element-wise."""
        z = np.clip(x.data, -500, 500) # Prevent overflow
        result_data = np.zeros_like(z)

        # Positive values: 1 / (1 + exp(-x))
        pos_mask = z >= 0
        result_data[pos_mask] = 1.0 / (1.0 + np.exp(-z[pos_mask]))

        # Negative values: exp(x) / (1 + exp(x))
```

```

neg_mask = z < 0
exp_z = np.exp(z[neg_mask])
result_data[neg_mask] = exp_z / (1.0 + exp_z)

return Tensor(result_data)

```

: Listing 2.1 — Numerically stable Sigmoid using a piecewise formula to avoid overflow on either side of zero. {#lst-02-activations-sigmoid}

Notice that there are two formulas, not one. The textbook form  $1 / (1 + \exp(-x))$  blows up for large negative  $x$  because  $\exp(-x)$  overflows. The algebraically equivalent  $\exp(x) / (1 + \exp(x))$  blows up for large positive  $x$  for the symmetric reason. Neither is “right”; each is right on one side of zero. The implementation picks the safe branch per element and clips at  $\pm 500$  as a belt-and-suspenders guard. This is the first time in the course you see a recurring pattern: *the math is one expression, the numerically stable code is a piecewise rewrite of it.*

Sigmoid’s smooth S-curve makes it natural for binary classification outputs — the value reads directly as a probability. For hidden layers it’s a disaster. When  $|x|$  grows, the output saturates near 0 or 1 and the gradient collapses toward zero. In a deep network those tiny gradients multiply through the chain rule and vanish exponentially. That’s why sigmoid lost the hidden-layer slot to ReLU around 2012 and never got it back.

Tanh is sigmoid’s zero-centered cousin, mapping inputs to  $(-1, 1)$ :

```

class Tanh:
    def forward(self, x: Tensor) -> Tensor:
        """Apply tanh activation element-wise."""
        result = np.tanh(x.data)
        return Tensor(result)

```

Zero-centering is the only real difference, and it matters more than it sounds. Sigmoid outputs are always positive, which biases the gradient updates of the next layer in one direction; tanh outputs straddle zero, so positive and negative gradients cancel naturally. That’s why tanh, not sigmoid, was the activation of choice inside LSTM and GRU cells where the same weights see the same hidden state hundreds of times. Tanh still saturates at the extremes, so deep stacks still vanish — but for bounded, recurrent settings it remains the right pick.

### 6.5.4 GELU

GELU (Gaussian Error Linear Unit) applies a more complex gate than ReLU. Rather than taking the  $\max(0, x)$  threshold, GELU applies a smooth, probabilistic transformation.

The implementation uses the fast sigmoid approximation from your module.

```

class GELU:
    def forward(self, x: Tensor) -> Tensor:
        """Apply GELU activation using the fast sigmoid approximation."""
        return Sigmoid()(x * 1.702) * x

```

: Listing 2.2 — GELU using the fast sigmoid approximation, reusing the Sigmoid implementation from this module. {#lst-02-activations-gelu}

While ReLU remains the default in CNNs and MLPs, GELU has become the standard choice for modern transformers like GPT and BERT.

**The tradeoff:** because it relies on transcendental operations (exponentials and divisions) rather than simple comparisons, it is roughly 4–5x more computationally expensive per element than ReLU. In the context of a billion-parameter model, this “GELU tax” is a deliberate choice: we trade raw throughput for the training stability and improved accuracy that a smooth activation provides at scale.

### 6.5.5 Softmax and Numerical Stability

Softmax converts any vector into a valid probability distribution. All outputs are positive, and they sum to exactly 1. This makes it essential for multi-class classification:

The code in `lst-02-activations-softmax` makes this concrete.

```
class Softmax:
    def forward(self, x: Tensor, dim: int = -1) -> Tensor:
        """Apply softmax activation along specified dimension."""
        # Numerical stability: subtract max to prevent overflow
        x_max_data = np.max(x.data, axis=dim, keepdims=True)
        x_max = Tensor(x_max_data, requires_grad=False)
        x_shifted = x - x_max

        # Compute exponentials
        exp_values = Tensor(np.exp(x_shifted.data),
                             requires_grad=x_shifted.requires_grad)

        # Sum along dimension
        exp_sum_data = np.sum(exp_values.data, axis=dim, keepdims=True)
        exp_sum = Tensor(exp_sum_data, requires_grad=exp_values.requires_grad)

        # Normalize to get probabilities
        result = exp_values / exp_sum
        return result
```

: Listing 2.3 — Softmax with max-subtraction for numerical stability (the log-sum-exp trick). (`lst-02-activations-softmax`)

**Why the max subtraction matters — the load-bearing trick of this chapter.**

`np.max` looks like a throwaway line. It is not. It is the single line that separates a softmax that works on real model logits from one that returns `nan` and silently corrupts your training run.

Consider what happens without it. A trained transformer can easily produce a logit of 50 for the predicted token.  $\exp(50)$  is about  $5.18 \times 10^{21}$  — still inside float32 range, but only barely. Push it to logit 100 and  $\exp(100) \approx 2.7 \times 10^{43}$  overflows to `inf`. Now the numerator is `inf`, the denominator is `inf`, and `inf / inf = nan`. Every downstream loss, gradient, and parameter update inherits the `nan`. By the time you notice, your model is ruined and the stack trace points at the loss function, not at the activation that poisoned it three layers earlier.

Subtracting the max fixes this without changing the math. The largest shifted logit is exactly 0, so the largest exponent we ever take is  $\exp(0) = 1$ . Every other exponent is between 0 and 1. Overflow becomes structurally impossible:

$$\begin{aligned} \text{softmax}(x)_i &= \exp(x_i - \max(x)) / \sum_j \exp(x_j - \max(x)) \\ &= \exp(x_i) / \exp(\max(x)) \div \sum_j \exp(x_j) / \exp(\max(x)) \\ &= \exp(x_i) / \sum_j \exp(x_j) \end{aligned}$$

The  $\exp(\max(x))$  factor cancels in numerator and denominator. Mathematically identical, numerically a different universe. This is the first instance of the **log-sum-exp trick** — you’ll meet it again in Module 04

when you implement cross-entropy loss as `log_softmax`, where the same shift prevents the *underflow* that happens when you take the log of a tiny probability. *That* is why every production framework — PyTorch, JAX, TensorFlow — does the same shift in the same place.

Two more properties worth internalizing before you move on.

**Softmax amplifies differences.** Input `[1, 2, 3]` becomes roughly `[0.09, 0.24, 0.67]`. The largest input is only  $3\times$  the smallest, but it claims 67% of the probability mass — because exponentials grow superlinearly. Push the inputs apart and the distribution sharpens toward one-hot; pull them together and it flattens toward uniform. This sharpening is what makes a classifier “confident.”

**Softmax couples its outputs.** Change one input and *every* output changes, because they all share the same denominator. That coupling is why the softmax gradient is a Jacobian, not an element-wise derivative — a complication you’ll inherit in Module 06 (Autograd) and the reason cross-entropy is fused with softmax in practice rather than backproped through both stages separately.

### 6.5.6 Choosing Activations

Here’s the decision tree production ML engineers actually use.

**For hidden layers:**

- Default: **ReLU** — fast, no vanishing gradients, creates sparsity
- Transformers: **GELU** — smooth, better gradient flow, the de-facto choice in GPT/BERT
- Recurrent networks: **Tanh** — zero-centered, plays well with repeated weight application
- When ReLU is dying on you: **LeakyReLU**, **ELU**, **Swish**

**For output layers:**

- Binary classification: **Sigmoid** — one probability in `[0, 1]`
- Multi-class classification: **Softmax** — a probability distribution that sums to 1
- Regression: **None** — leave the linear output alone

**Computational cost (relative to ReLU):**

- ReLU:  $1\times$  (just a comparison)
- Sigmoid/Tanh:  $3\text{--}4\times$  (one exponential per element)
- GELU:  $4\text{--}5\times$  (exponential plus the approximation polynomial)
- Softmax:  $5\times+$  (exponential, sum-reduction, division)

For a billion-parameter model, swapping ReLU for GELU in every hidden layer can add 20–30% to training time. The trade is often worth it: a 1–2 point accuracy bump on a transformer is rarely something you turn down. The point is that you have to *know it’s a trade*. Defaults shipped by frameworks are not free.

### 6.5.7 Computational Complexity

All activation functions are element-wise operations, meaning they apply independently to each element of the tensor. This gives  $O(n)$  time complexity where  $n$  is the total number of elements. However, because they perform very little math per element, activation functions are typically **memory-bound** rather than compute-bound. The bottleneck isn’t the processor doing the math; it’s the time it takes to fetch the tensor from RAM and write the result back.

This physical constraint fundamentally changes how systems engineers implement these seemingly simple mathematical functions. When compute is fast but memory access is slow, every byte moved across the bus is a tax on performance.

### **i Systems Implication: The Bandwidth Tax & In-place Operations**

Because element-wise operations like ReLU spend most of their execution time waiting on memory transfers, production frameworks often use **in-place operations** (e.g., `x.relu_()`). By overwriting the input tensor directly rather than allocating a new one, in-place operations save an entire round-trip to RAM, significantly reducing the bandwidth tax.

Even though these operations are memory-bound, the constant mathematical factors still differ dramatically:

Table 6.3 compares their cost against the ReLU baseline.

Table 6.3: **Computational cost of activations relative to ReLU.**

| Operation             | Complexity                                          | Cost Relative to ReLU |
|-----------------------|-----------------------------------------------------|-----------------------|
| ReLU ( $\max(0, x)$ ) | $O(n)$ comparisons                                  | $1\times$ (baseline)  |
| Sigmoid/Tanh          | $O(n)$ exponentials                                 | $3\text{--}4\times$   |
| GELU                  | $O(n)$ exponentials + multiplies                    | $4\text{--}5\times$   |
| Softmax               | $O(n)$ exponentials + $O(n)$ sum + $O(n)$ divisions | $5\times+$            |

Exponentials are expensive. A modern CPU can execute 1 billion comparisons per second but only 250 million exponentials per second. This is why ReLU is so popular: at scale, a  $4\times$  speedup in activation computation can mean the difference between training in 1 day versus 4 days.

Memory complexity is  $O(n)$  for all activations because they create an output tensor the same size as the input. Softmax requires small temporary buffers for the exponentials and sum, but this overhead is negligible compared to the tensor sizes in production networks.

## 6.6 Production Context

### 6.6.1 Your Implementation vs. PyTorch

Your TinyTorch activations and PyTorch's `torch.nn.functional` activations implement the same mathematical functions with the same numerical stability measures. The differences are in optimization and GPU support:

Table 6.4 places your implementation side by side with the production reference for direct comparison.

Table 6.4: **Feature comparison between TinyTorch activations and PyTorch equivalents.**

| Feature                    | Your Implementation                           | PyTorch                                     |
|----------------------------|-----------------------------------------------|---------------------------------------------|
| <b>Backend</b>             | NumPy (Python/C)                              | C++/CUDA kernels                            |
| <b>Speed</b>               | $1\times$ (CPU baseline)                      | $10\text{--}100\times$ faster (GPU)         |
| <b>Numerical Stability</b> | Max subtraction (Softmax), clipping (Sigmoid) | Same techniques                             |
| <b>Autograd</b>            | Stubs (added later)                           | Full gradient computation                   |
| <b>Variants</b>            | 5 core activations                            | 30+ variants (LeakyReLU, PReLU, Mish, etc.) |

### 6.6.2 Code Comparison

The following comparison shows equivalent activation usage in TinyTorch and PyTorch. Notice how the APIs are nearly identical, differing only in import paths and minor syntax.

## 6.7 Your TinyTorch

```
from tinytorch.core.activations import ReLU, Sigmoid, Softmax
from tinytorch.core.tensor import Tensor

# Element-wise activations
x = Tensor([[ -1, 0, 1, 2]])
relu = ReLU()
activated = relu(x) # [0, 0, 1, 2]

# Binary classification output
sigmoid = Sigmoid()
probability = sigmoid(x) # All values in (0, 1)

# Multi-class classification output
logits = Tensor([[1, 2, 3]])
softmax = Softmax()
probs = softmax(logits) # [0.09, 0.24, 0.67], sum = 1
```

## 6.8 PyTorch

```
import torch
import torch.nn.functional as F

# Element-wise activations
x = torch.tensor([[ -1, 0, 1, 2]], dtype=torch.float32)
activated = F.relu(x) # [0, 0, 1, 2]

# Binary classification output
probability = torch.sigmoid(x) # All values in (0, 1)

# Multi-class classification output
logits = torch.tensor([[1, 2, 3]], dtype=torch.float32)
probs = F.softmax(logits, dim=-1) # [0.09, 0.24, 0.67], sum = 1
```

Let's walk through the key similarities and differences:

- **Line 1 (Import):** TinyTorch imports activation classes; PyTorch uses functional interface `torch.nn.functional`. Both approaches work; PyTorch also supports class-based activations via `torch.nn.ReLU()`.
- **Line 4-6 (ReLU):** Identical semantics. Both zero out negative values, preserving positive values.
- **Line 9-10 (Sigmoid):** Identical mathematical function. Both use numerically stable implementations to prevent overflow.

- **Line 13-15 (Softmax):** Same mathematical operation. Both require specifying the dimension for multidimensional tensors. PyTorch uses `dim` keyword argument; TinyTorch defaults to `dim=-1`.

#### 💡 What's Identical

Mathematical functions, numerical stability techniques (max subtraction in softmax), and the concept of element-wise transformations. When you debug PyTorch activation issues, you'll understand exactly what's happening because you implemented the same logic.

### 6.8.1 Why Activations Matter at Scale

A single activation call is microseconds of work. The reason it adds up to a real engineering trade-off is the multiplier.

- **Large language models:** GPT-3 has 96 transformer layers, each with 2 GELU activations — **192 GELU operations per forward pass**, every one of them over a tensor with billions of elements.
- **Image classification:** ResNet-50 has 49 convolutional layers, each followed by ReLU. A batch of 256 images at  $224 \times 224$  resolution executes roughly **12 billion ReLU operations** per batch.
- **Production serving:** A model handling 1000 requests per second runs about **86 million activation computations per day**, per layer. A 20% speedup from picking ReLU over GELU translates to hours of saved compute every day.

Activations account for **5–15% of total training time** in a typical conv-heavy network. In transformers — where the per-layer matmuls are smaller and the layer count is larger — that share climbs to **20–30%**. That's the cost side of the GELU-vs-ReLU choice: not negligible, not catastrophic, but absolutely worth measuring before you ship.

## 6.9 Check Your Understanding

### 💡 Check Your Understanding — Activations

Before moving on, verify you can articulate each of the following:

- Why stacking linear layers without activations collapses to a single matrix multiply — and how ReLU, Sigmoid, or GELU break that collapse.
- Why softmax subtracts the per-row max before exponentiating, and how that one shift turns an `inf/inf = nan` blowup into a stable `exp(0) = 1`.
- The dying ReLU failure chain (negative pre-activation → zero output → zero local gradient → frozen weights) and why LeakyReLU/GELU patch it.
- Why activations are memory-bound (not compute-bound) and how in-place variants like `relu_()` save an entire RAM round-trip.
- The relative cost ordering  $\text{ReLU} (1x) < \text{Sigmoid/Tanh} (3\text{--}4x) < \text{GELU} (4\text{--}5x) < \text{Softmax} (5x+)$ , and when that cost is worth paying.

If any of these feels fuzzy, revisit Core Concepts (Why Non-linearity Matters, Softmax and Numerical Stability, Computational Complexity) before moving on.

The collapsible Q&A below works each of these through with numbers.

#### Q1: Memory Calculation

A batch of 32 samples passes through a hidden layer with 4096 neurons and ReLU activation. How much memory is required to store the activation outputs (float32)?

 **Answer**

$32 \times 4096 \times 4 \text{ bytes} = 524,288 \text{ bytes} \approx 512 \text{ KB}$ .

This is the activation memory for ONE layer. A 100-layer network needs **50 MB** just to store activations for a single forward pass — and during training every one of those layers has to be cached so backprop can use them. That's why activation memory, not parameter memory, is what usually decides your batch size.

**Q2: Computational Cost**

If ReLU takes 1ms to activate 1 million neurons, approximately how long will GELU take on the same input?

 **Answer**

GELU is approximately **4–5× slower** than ReLU due to exponential computation in the sigmoid approximation.

Expected time: **4–5ms**

At scale, this matters: if you have 100 activation layers in your model, switching from ReLU to GELU adds 300–400ms per forward pass. For training that requires millions of forward passes, this multiplies into hours or days of extra compute time.

**Q3: Numerical Stability**

Why does softmax subtract the maximum value before computing exponentials? What would happen without this step?

 **Answer**

**Without max subtraction:** Computing `softmax([1000, 1001, 1002])` requires  $\exp(1000)$ , which overflows to infinity in float32/float64, producing NaN.

**With max subtraction:** First compute  $x_{\text{shifted}} = x - \max(x) = [0, 1, 2]$ , then compute  $\exp([0, 1, 2])$  which stays within float range.

**Why this works mathematically:**

$$\begin{aligned} \exp(x - \max) / \sum \exp(x - \max) &= [\exp(x) / \exp(\max)] / [\sum \exp(x) / \exp(\max)] \\ &= \exp(x) / \sum \exp(x) \end{aligned}$$

The  $\exp(\max)$  factor cancels out, so the result is mathematically identical. But numerically, it prevents overflow. This is a classic example of why production ML requires careful numerical engineering, not just correct math.

**Q4: Sparsity Analysis**

A ReLU layer processes input tensor with shape (128, 1024) containing values drawn from a normal distribution  $N(0, 1)$ . Approximately what percentage of outputs will be exactly zero?

 **Answer**

For a standard normal distribution  $N(0, 1)$ , approximately **50% of values are negative**.

ReLU zeros all negative values, so approximately **50% of outputs will be exactly zero**.

Total elements:  $128 \times 1024 = 131,072$ . Zeros:  $\approx 65,536$ .

This sparsity has major implications:

- **Speed:** Multiplying by zero is free, so downstream computations can skip ~50% of operations
- **Memory:** Sparse formats can compress the output by  $2\times$



- **Generalization:** Sparse representations often generalize better (less overfitting)

This is why ReLU is so effective: it creates natural sparsity without requiring explicit regularization.

#### Q5: Activation Selection

You're building a sentiment classifier that outputs "positive" or "negative". Which activation should you use for the output layer, and why?

#### 💡 Answer

Use **Sigmoid** for the output layer.

**Reasoning:**

- Binary classification needs a single probability value in  $[0, 1]$
- Sigmoid maps any real number to  $(0, 1)$
- Output can be interpreted as  $P(\text{positive})$  where 0.8 means "80% confident this is positive"
- Decision rule: predict positive if  $\text{sigmoid}(\text{output}) > 0.5$

**Why NOT other activations:**

- **Softmax:** Overkill for binary classification (designed for multi-class), though technically works with 2 outputs
- **ReLU:** Outputs unbounded positive values, not interpretable as probabilities
- **Tanh:** Outputs in  $(-1, 1)$ , not directly interpretable as probabilities

**Production pattern:**

Input  $\rightarrow$  Linear + ReLU  $\rightarrow$  Linear + ReLU  $\rightarrow$  Linear + Sigmoid  $\rightarrow$  Binary Probability

For multi-class sentiment (positive/negative/neutral), you'd use Softmax instead to get a 3-element probability distribution.

## 6.10 Key Takeaways

- **Nonlinearity is non-optional:** without a pointwise function between layers, any depth of Linear stack collapses to a single  $wx$  — activations are what make "deep" mean anything.
- **The log-sum-exp shift is load-bearing:** every production softmax subtracts  $\max(x)$  before exponentiating, turning overflow into algebraically identical safe math. You will see this trick again inside cross-entropy.
- **ReLU trades elegance for dead neurons:**  $\max(0, x)$  is the cheapest possible nonlinearity and gives constant gradient on the positive side, but any weight update that zeros the pre-activation permanently silences that unit.
- **Activations are memory-bound:** per-element compute is trivial; the bottleneck is the RAM round-trip, which is why in-place `x.relu_()` is preferred in production and why the ReLU $\rightarrow$ GELU swap is a measurable cost.
- **Cost scales with transcendentals:** ReLU is one comparison; Sigmoid/Tanh need one `exp`; Softmax needs an `exp`, a reduction, and a division. That ordering decides when GELU's smoother gradient is worth its 4-5x tax.

**Coming next:** Module 03 adds the learnable transformation between activations — the `Linear` layer and the `Sequential` container that lets you compose `Linear  $\rightarrow$  activation  $\rightarrow$  Linear  $\rightarrow$  ...` into an actual MLP.

## 6.11 Further Reading

The choice of an activation function is rarely just a mathematical preference; it is a profound architectural decision that dictates how a network learns, sparsifies data, and how efficiently it maps to hardware. To understand how we arrived at today's standard practices—from the memory-efficient sparsity of ReLU to the computationally heavy but smooth gradients of GELU—the following historical foundations are essential reading.

### 6.11.1 Seminal Papers

- **Deep Sparse Rectifier Neural Networks** - Glorot, Bordes, Bengio (2011). The paper that established ReLU as the default activation for deep networks, showing how its sparsity and constant gradient enable training of very deep networks. [AISTATS](#)
- **Gaussian Error Linear Units (GELUs)** - Hendrycks & Gimpel (2016). Introduced the smooth activation that powers modern transformers like GPT and BERT. Explains the probabilistic interpretation and why smoothness helps optimization. [arXiv:1606.08415](#)
- **Attention Is All You Need** - Vaswani et al. (2017). While primarily about transformers, this paper's use of specific activations (ReLU in position-wise FFN, Softmax in attention) established patterns still used today. **Systems Implication:** Replaced recurrence with self-attention, breaking the sequential compute bottleneck of RNNs and allowing massive parallelization across GPU cores. [NeurIPS](#)

### 6.11.2 Additional Resources

- **Textbook:** “[Deep Learning](#)” by Goodfellow, Bengio, and Courville - Chapter 6.3 covers activation functions with mathematical rigor
- **Blog:** [Understanding Activation Functions](#) - Amazon's MLU visual explanation of ReLU

## 6.12 What's Next

You now have nonlinearity. You still don't have anything *to* be nonlinear about. ReLU on a raw input vector accomplishes nothing — the network needs a learnable transformation between activations, something with weights and biases that gradient descent can shape.

That's the next module.

### Coming Up: Module 03 — Layers

**The question Module 03 answers:** what is `Linear(x)`, exactly, and how does it compose with the activations you just built into the canonical `Linear → activation → Linear → activation → ...` stack that defines an MLP?

You'll implement the `Linear` layer — weight matrix, bias vector, forward pass — and then chain it with your `ReLU` and `Softmax` to build the first thing in this course that deserves to be called a neural network.

### Preview — how your activations get used in future modules:

Table 6.5 traces how this module is reused by later parts of the curriculum.

Table 6.5: How activations feed into subsequent TinyTorch modules.

| Module       | What It Does                   | Your Activations In Action                                      |
|--------------|--------------------------------|-----------------------------------------------------------------|
| 03: Layers   | Neural network building blocks | <code>Linear(x)</code> followed by <code>ReLU()</code> (output) |
| 04: Losses   | Training objectives            | Softmax probabilities feed into cross-entropy loss              |
| 06: Autograd | Automatic gradients            | <code>relu.backward(grad)</code> computes activation gradients  |

## 6.13 Get Started

### 💡 Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### ⚠️ Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



 Chapter 7

## Module 03: Layers

A layer is a function with weights. Three stacked is an MLP; ninety-six is GPT-3. The interface discipline, every layer exposing the same `forward()` and `parameters()`, is what lets the optimizer find your weights, autograd walk your graph, and a future compiler fuse your kernels.

### Module Info

**FOUNDATION TIER** | Difficulty: ●●○○ | Time: 5-7 hours | Prerequisites: 01, 02

**Prerequisites: Modules 01 and 02** means you have built:

- Tensor class with arithmetic, broadcasting, matrix multiplication, and shape manipulation
- Activation functions (ReLU, Sigmoid, Tanh, Softmax) for introducing non-linearity
- Understanding of element-wise operations and reductions

If you can multiply tensors, apply activations, and reason about shape transformations, you're ready.

### 7.1 Overview

A layer is a function with weights. Stack three of them and you have a multi-layer perceptron; stack ninety-six and you have GPT-3. The discipline that makes such composition possible — every layer exposing the same `forward()` and `parameters()` interface — is what you build in this module.

You'll implement four pieces: a `Layer` base class that fixes the interface, a `Linear` layer that applies the learned transformation  $y = xW + b$ , a `Dropout` layer that prevents overfitting by randomly zeroing activations, and a `Sequential` container that chains layers into networks. Together they're the smallest set of abstractions that lets you write `model = Sequential(Linear(784, 256), ReLU(), Linear(256, 10))` and have it just work.

The payoff is `parameters()`. Once every layer hands its weights to a single list, optimizers (Module 07) and autograd (Module 06) can update them without caring what's inside. PyTorch's `nn.Module` follows the exact same contract — you're building the real abstraction, not a toy version of it.

### 7.2 Learning Objectives

#### By completing this module, you will:

- **Implement** Linear layers with proper weight initialization and parameter management for gradient-based training
- **Master** the mathematical operation  $y = xW + b$  and understand how parameter counts scale with layer dimensions
- **Understand** memory usage patterns (parameter memory vs activation memory) and computational complexity of matrix operations

- **Connect** your implementation to production PyTorch patterns, including `nn.Linear`, `nn.Dropout`, and parameter tracking

## 7.3 What You'll Build




Figure 7.1: **One interface, three subclasses.** Linear, Dropout, and Sequential all inherit `forward()` and `parameters()` from a single Layer base — Sequential adds composition by chaining the others into networks.

### Implementation roadmap:

Table 7.1 lays out the implementation in order, one part at a time.

Table 7.1: **Implementation roadmap for the Layer, Linear, Dropout, and Sequential classes.**

| Part | What You'll Implement                                                                              | Key Concept                           |
|------|----------------------------------------------------------------------------------------------------|---------------------------------------|
| 1    | Layer base class with <code>forward()</code> , <code>__call__()</code> , <code>parameters()</code> | Consistent interface for all layers   |
| 2    | Linear layer with proper initialization                                                            | Learned transformation $y = xW + b$   |
| 3    | Dropout with training/inference modes                                                              | Regularization through random masking |
| 4    | Sequential container for layer composition                                                         | Chaining layers together              |

### The pattern you'll enable:

```

# Building a multi-layer network
layer1 = Linear(784, 256)
activation = ReLU()
dropout = Dropout(0.5)
layer2 = Linear(256, 10)

# Manual composition for explicit data flow
x = layer1(x)
x = activation(x)
x = dropout(x, training=True)

```

```
output = layer2(x)
```

### 7.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Automatic gradient computation (autograd is a later module)
- Parameter optimization (optimizers are a later module)
- Hundreds of layer types (PyTorch has Conv2d, LSTM, Attention - you'll build Linear and Dropout)
- Automatic training/eval mode switching (PyTorch's `model.train()` - you'll manually pass `training` flag)

**You are building the core building blocks.** Training loops and optimizers come later.

## 7.4 API Reference

A quick reference for the layer classes you're about to build — keep it open while you implement.

### 7.4.1 Layer Base Class

```
Layer()
```

Base class providing consistent interface for all neural network layers. All layers inherit from this and implement `forward()` and `parameters()`.

Table 7.2 lists the methods every subclass inherits.

Table 7.2: Methods defined by the Layer base class.

| Method                  | Signature                                    | Description                          |
|-------------------------|----------------------------------------------|--------------------------------------|
| <code>forward</code>    | <code>forward(x) -&gt; Tensor</code>         | Compute layer output (must override) |
| <code>__call__</code>   | <code>__call__(x) -&gt; Tensor</code>        | Makes layer callable like a function |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code> | Returns list of trainable parameters |

### 7.4.2 Linear Layer

```
Linear(in_features, out_features, bias=True)
```

Linear (fully connected) layer implementing  $y = xW + b$ .

**Parameters:** - `in_features`: Number of input features - `out_features`: Number of output features - `bias`: Whether to include bias term (default: `True`)

**Attributes:** - `weight`: Tensor of shape `(in_features, out_features)` (gradient tracking enabled later by autograd) - `bias`: Tensor of shape `(out_features,)` or `None` (gradient tracking enabled later by autograd)

Table 7.3 lists the methods on this class.

Table 7.3: **Methods on the Linear layer.**

| Method                  | Signature                                    | Description                                                  |
|-------------------------|----------------------------------------------|--------------------------------------------------------------|
| <code>forward</code>    | <code>forward(x) -&gt; Tensor</code>         | Applies linear transformation $y = xW + b$                   |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code> | Returns <code>[weight, bias]</code> or <code>[weight]</code> |

### 7.4.3 Dropout Layer

```
Dropout(p=0.5)
```

Dropout layer for regularization. During training, randomly zeros elements with probability  $p$  and scales survivors by  $1/(1-p)$ . During inference, passes input unchanged.

**Parameters:** -  $p$ : Probability of zeroing each element (0.0 = no dropout, 1.0 = zero everything)

Table 7.4 lists the methods on this class.

Table 7.4: **Methods on the Dropout layer.**

| Method                  | Signature                                           | Description                                                   |
|-------------------------|-----------------------------------------------------|---------------------------------------------------------------|
| <code>forward</code>    | <code>forward(x, training=True) -&gt; Tensor</code> | Applies dropout during training, passthrough during inference |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code>        | Returns empty list (no trainable parameters)                  |

### 7.4.4 Sequential Container

```
Sequential(*layers)
```

Container that chains layers together sequentially. Provides convenient way to compose multiple layers. Table 7.5 lists the methods on this class.

Table 7.5: **Methods on the Sequential container.**

| Method                  | Signature                                    | Description                              |
|-------------------------|----------------------------------------------|------------------------------------------|
| <code>forward</code>    | <code>forward(x) -&gt; Tensor</code>         | Forward pass through all layers in order |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code> | Collects all parameters from all layers  |

## 7.5 Core Concepts

Five ideas carry the weight of this module: the linear transformation, weight initialization, parameter tracking, the forward-pass call mechanic, and composition. They apply to every ML framework. TinyTorch is the lab; the concepts are what travel.



### 7.5.1 The Linear Transformation

Linear layers implement the mathematical operation  $y = xW + b$ , where  $x$  is your input,  $W$  is a weight matrix you learn,  $b$  is a bias vector you learn, and  $y$  is your output. This simple formula is the foundation of neural networks.

Think of the weight matrix as a feature detector. Each column of  $W$  learns to recognize a particular pattern in the input. When you multiply input  $x$  by  $W$ , you're asking: "How much of each learned pattern appears in this input?" The bias  $b$  shifts the output, providing a baseline independent of the input.

Consider recognizing handwritten digits. A flattened  $28 \times 28$  image has 784 pixels. A Linear layer transforming 784 features to 10 classes creates a weight matrix of shape  $(784, 10)$ . Each of the 10 columns learns which combination of those 784 pixels indicates a particular digit. The network discovers these patterns through training.

Here's how your implementation performs this transformation:

```
def forward(self, x):  
    """Forward pass through linear layer."""  
    # Linear transformation:  $y = xW$   
    output = x.matmul(self.weight)  
  
    # Add bias if present  
    if self.bias is not None:  
        output = output + self.bias  
  
    return output
```

The elegance is in the simplicity, but beneath that simplicity lies a massive computational engine. This matrix multiplication—formally known as **GEMM** (General Matrix Multiply)—handles all the feature combinations in one operation. Unlike memory-bound activation functions, GEMM has incredibly high **arithmetic intensity** (the ratio of math operations to memory bytes accessed). Because each fetched weight and input is reused multiple times across the multiplication, the processor can spend its time crunching numbers rather than waiting for RAM.

However, realizing this theoretical performance requires moving beyond textbook linear algebra. If we just wrote a naive nested loop for matrix multiplication, the processor would spend most of its time stalling for memory fetches, completely wasting its computational potential. Bridging this gap requires specialized software-hardware co-design.

#### **i** Systems Implication: GEMM & Cache Tiling

To maximize the benefits of high arithmetic intensity, modern hardware relies on **cache tiling**. Instead of naively multiplying rows and columns (which thrashes the CPU/GPU cache), hardware libraries break the matrices into smaller "tiles" that fit perfectly into ultra-fast L1/L2 caches. The math happens on these tiles at blazing speeds before moving to the next block, ensuring the execution remains compute-bound rather than memory-bound.

After the heavy lifting of GEMM is done, broadcasting efficiently handles adding the bias vector to every sample in the batch. This single `matmul` method essentially powers every linear transformation in modern neural networks.

### 7.5.2 Weight Initialization

How you initialize weights determines whether your network can learn at all. Initialize too small and gradients vanish, making learning impossibly slow. Initialize too large and gradients explode, making training unstable. The sweet spot ensures stable gradient flow through the network.

We use LeCun-style initialization, which scales weights by `sqrt(1/in_features)`. This keeps the variance of activations roughly constant as data flows through layers, preventing vanishing or exploding gradients. (True Xavier/Glorot uses `sqrt(2/(fan_in+fan_out))` which also considers output dimensions, but the simpler LeCun formula works well in practice.)

Here's your initialization code:

The code in `?lst-03-layers-linear-init` makes this concrete.

```
def __init__(self, in_features, out_features, bias=True):
    """Initialize linear layer with proper weight initialization."""
    self.in_features = in_features
    self.out_features = out_features

    # LeCun-style initialization for stable gradients
    scale = np.sqrt(INIT_SCALE_FACTOR / in_features)
    weight_data = np.random.randn(in_features, out_features) * scale
    self.weight = Tensor(weight_data)

    # Initialize bias to zeros or None
    if bias:
        bias_data = np.zeros(out_features)
        self.bias = Tensor(bias_data)
    else:
        self.bias = None
```

: Listing 3.1 — Linear layer constructor with LeCun-style weight initialization and zero bias. {#lst-03-layers-linear-init}

Weights and biases are created as plain Tensors. When autograd is enabled later, it monkey-patches the Tensor class to support `requires_grad`, at which point you can set `layer.weight.requires_grad = True` for parameters that need gradients. Bias starts at zero because the weight initialization already handles the scale, and zero is a neutral starting point for per-class adjustments.

For `Linear(1000, 10)`, the scale is `sqrt(1/1000) ≈ 0.032`. For `Linear(10, 1000)`, the scale is `sqrt(1/10) ≈ 0.316`. Layers with more inputs get smaller initial weights because each input contributes to the output, and you want their combined effect to remain stable.

### 7.5.3 Parameter Management

Parameters are tensors that need gradients and optimizer updates. Your Linear layer manages two parameters: weights and biases. The `parameters()` method collects them into a list that optimizers can iterate over.

```
def parameters(self):
    """Return list of trainable parameters."""
    params = [self.weight]
    if self.bias is not None:
        params.append(self.bias)
```

```
return params
```

This simple method enables powerful workflows. When you build a multi-layer network, you can collect all parameters from all layers and pass them to an optimizer:

```
layer1 = Linear(784, 256)
layer2 = Linear(256, 10)

all_params = layer1.parameters() + layer2.parameters()
# Pass all_params to optimizer.step() during training
```

Each Linear layer independently manages its own parameters. The Sequential container extends this pattern by collecting parameters from all its contained layers, enabling hierarchical composition.

#### 7.5.4 Forward Pass Mechanics

The forward pass transforms input data through the layer's computation. Every layer implements `forward()`, and the base class provides `__call__()` to make layers callable like functions. This matches PyTorch's design exactly.

```
def __call__(self, x, *args, **kwargs):
    """Allow layer to be called like a function."""
    return self.forward(x, *args, **kwargs)
```

This lets you write `output = layer(input)` instead of `output = layer.forward(input)`. The difference looks cosmetic, but `__call__` is the seam where a real framework slips in hooks, logging, profiling, or training/eval mode switching — none of which `forward()` has to know about.

For Dropout, the forward pass depends on whether you're training or performing inference:

The code in `?lst-03-layers-dropout-forward` makes this concrete.

```
def forward(self, x, training=True):
    """Forward pass through dropout layer."""
    if not training or self.p == DROPOUT_MIN_PROB:
        # During inference or no dropout, pass through unchanged
        return x

    if self.p == DROPOUT_MAX_PROB:
        # Drop everything (preserve requires_grad for gradient flow)
        return Tensor(np.zeros_like(x.data), requires_grad=x.requires_grad)

    # During training, apply dropout
    keep_prob = 1.0 - self.p

    # Create random mask: True where we keep elements
    mask = np.random.random(x.data.shape) < keep_prob

    # Apply mask and scale using Tensor operations to preserve gradients
    mask_tensor = Tensor(mask.astype(np.float32), requires_grad=False)
    scale = Tensor(np.array(1.0 / keep_prob), requires_grad=False)
```

```
# Use Tensor operations: x * mask * scale
output = x * mask_tensor * scale
return output
```

: Listing 3.2 — Dropout forward pass with inverted scaling  $1/(1-p)$  to preserve expected magnitude. {#lst-03-layers-dropout-forward}

The key insight is the scaling factor  $1/(1-p)$ . If you drop 50% of neurons, the survivors need to be scaled by 2.0 to maintain the same expected value. This ensures that during inference (when no dropout is applied), the output magnitudes match training expectations.

### 7.5.5 Layer Composition

Neural networks are built by chaining layers together. Data flows through each layer in sequence, with each transformation building on the previous one. Your Sequential container captures this pattern:

The code in ?@lst-03-layers-sequential makes this concrete.

```
class Sequential:
    """Container that chains layers together sequentially."""

    def __init__(self, *layers):
        """Initialize with layers to chain together."""
        if len(layers) == 1 and isinstance(layers[0], (list, tuple)):
            self.layers = list(layers[0])
        else:
            self.layers = list(layers)

    def forward(self, x):
        """Forward pass through all layers sequentially."""
        for layer in self.layers:
            x = layer.forward(x)
        return x

    def parameters(self):
        """Collect all parameters from all layers."""
        params = []
        for layer in self.layers:
            params.extend(layer.parameters())
        return params
```

: Listing 3.3 — Sequential container chaining forward passes and flattening per-layer parameters. {#lst-03-layers-sequential}

Twelve lines of container code unlock the full vocabulary of network design. A 3-layer classifier is just three Linear layers with activations and dropout between them:

```
model = Sequential(
    Linear(784, 256), ReLU(), Dropout(0.5),
    Linear(256, 128), ReLU(), Dropout(0.3),
    Linear(128, 10)
```

)

`forward()` chains the computations; `parameters()` flattens the per-layer lists into one. That's the entire mechanism — every model in this book, from a perceptron to a transformer, reduces to those two methods on a container.

### 7.5.6 Memory and Computational Complexity

Memory and FLOPs aren't optional reading — they're how you predict whether a model fits in RAM and whether a forward pass takes a millisecond or a minute. Linear layers dominate both for any fully connected architecture.

Parameter memory for a Linear layer is straightforward:  $\text{in\_features} \times \text{out\_features} \times 4$  bytes for weights, plus  $\text{out\_features} \times 4$  bytes for bias (assuming float32). For `Linear(784, 256)`:

```
Weights: 784 × 256 × 4 = 802,816 bytes ≈ 784 KB
Bias:    256 × 4      =   1,024 bytes ≈   1 KB
Total:                   ≈ 785 KB
```

Activation memory depends on batch size. For batch size 32 and the same layer:

```
Input:   32 × 784 × 4 = 100,352 bytes ≈ 98 KB
Output:  32 × 256 × 4 =  32,768 bytes ≈ 32 KB
```

The computational cost of the forward pass is dominated by matrix multiplication. For input shape  $(\text{batch}, \text{in\_features})$  and weight shape  $(\text{in\_features}, \text{out\_features})$ , the operation requires  $\text{batch} \times \text{in\_features} \times \text{out\_features}$  multiplications and the same number of additions. Bias addition is just  $\text{batch} \times \text{out\_features}$  additions, negligible compared to matrix multiplication.

Table 7.6 summarises the complexity and speed of each stage.

Table 7.6: Compute and memory cost of Linear and Dropout forward passes.

| Operation         | Complexity                                           | Memory                                                           |
|-------------------|------------------------------------------------------|------------------------------------------------------------------|
| Linear forward    | $O(\text{batch} \times \text{in} \times \text{out})$ | $O(\text{batch} \times (\text{in} + \text{out}))$<br>activations |
| Dropout forward   | $O(\text{batch} \times \text{features})$             | $O(\text{batch} \times \text{features})$ mask                    |
| Parameter storage | $O(\text{in} \times \text{out})$                     | $O(\text{in} \times \text{out})$ weights                         |

For a 3-layer network (784→256→128→10) with batch size 32:

```
Layer 1: 32 × 784 × 256 = 6,422,528 FLOPs
Layer 2: 32 × 256 × 128 = 1,048,576 FLOPs
Layer 3: 32 × 128 ×  10 =    40,960 FLOPs
Total:                   ≈ 7.5 million FLOPs per forward pass
```

The first layer dominates because it has the largest input dimension. Production networks often shrink dimensions early to save computation in deeper layers.

## 7.6 Common Errors

These are the errors you'll encounter most often when working with layers. Understanding why they happen will save you hours of debugging, both in this module and throughout your ML career.

### 7.6.1 Shape Mismatch in Layer Composition

**Error:** `ValueError: Cannot perform matrix multiplication: (32, 128) @ (256, 10). Inner dimensions must match: 128  $\neq$  256`

This happens when you chain layers with incompatible dimensions. If `layer1` outputs 128 features but `layer2` expects 256 input features, the matrix multiplication in `layer2` fails.

**Fix:** Ensure output features of one layer match input features of the next:

```
layer1 = Linear(784, 128) # Outputs 128 features
layer2 = Linear(128, 10)  # Expects 128 input features
```

### 7.6.2 Dropout in Inference Mode

**Error:** Test accuracy is much lower than training accuracy, but loss curves suggest good learning

**Cause:** You're applying dropout during inference. Dropout should only zero elements during training. During inference, all neurons must be active.

**Fix:** Always pass `training=False` during evaluation:

```
# Training
output = dropout(x, training=True)

# Evaluation
output = dropout(x, training=False)
```

### 7.6.3 Missing Parameters

**Error:** Optimizer has no parameters to update, or parameter count is wrong

**Cause:** Your `parameters()` method doesn't return all trainable tensors, or you forgot to set `requires_grad=True`.

**Fix:** Verify all tensors with `requires_grad=True` are returned:

```
def parameters(self):
    params = [self.weight]
    if self.bias is not None:
        params.append(self.bias)
    return params # Must include all trainable tensors
```

### 7.6.4 Initialization Scale

**Error:** Loss becomes NaN within a few iterations, or gradients vanish immediately

**Cause:** Weights initialized too large (exploding gradients) or too small (vanishing gradients).

**Fix:** Use proper initialization scaling:

```
scale = np.sqrt(1.0 / in_features)  # LeCun-style, not just random()!
weight_data = np.random.randn(in_features, out_features) * scale
```

## 7.7 Production Context

### 7.7.1 Your Implementation vs. PyTorch

Your TinyTorch layers and PyTorch's `nn.Linear` and `nn.Dropout` share the same conceptual design. The differences are in implementation details: PyTorch uses C++ for speed, supports GPU acceleration, and provides hundreds of specialized layer types. But the core abstractions are identical.

Table 7.7 places your implementation side by side with the production reference for direct comparison.

Table 7.7: Feature comparison between TinyTorch layers and PyTorch's `nn` module.

| Feature                     | Your Implementation                   | PyTorch                                                                                          |
|-----------------------------|---------------------------------------|--------------------------------------------------------------------------------------------------|
| <b>Backend</b>              | NumPy (Python)                        | C++/CUDA                                                                                         |
| <b>Initialization</b>       | LeCun-style manual                    | Multiple schemes<br>( <code>init.xavier_uniform_</code> ,<br><code>init.kaiming_normal_</code> ) |
| <b>Parameter Management</b> | Manual <code>parameters()</code> list | <code>nn.Module</code> base class<br>with auto-registration                                      |
| <b>Training Mode</b>        | Manual <code>training</code> flag     | <code>model.train()</code> /<br><code>model.eval()</code> state                                  |
| <b>Layer Types</b>          | Linear, Dropout                       | 100+ layer types<br>(Conv, LSTM,<br>Attention, etc.)                                             |
| <b>GPU Support</b>          | CPU only                              | CUDA, Metal, ROCm                                                                                |

### 7.7.2 Code Comparison

The following comparison shows equivalent layer operations in TinyTorch and PyTorch. Notice how closely the APIs mirror each other.

## 7.8 Your TinyTorch

```
from tinytorch.core.layers import Linear, Dropout, Sequential
from tinytorch.core.activations import ReLU

# Build layers
layer1 = Linear(784, 256)
activation = ReLU()
dropout = Dropout(0.5)
layer2 = Linear(256, 10)

# Manual composition
x = layer1(x)
```

```

x = activation(x)
x = dropout(x, training=True)
output = layer2(x)

# Or use Sequential
model = Sequential(
    Linear(784, 256), ReLU(), Dropout(0.5),
    Linear(256, 10)
)
output = model(x)

# Collect parameters
params = model.parameters()

```

## 7.9 PyTorch

```

import torch
import torch.nn as nn

# Build layers
layer1 = nn.Linear(784, 256)
activation = nn.ReLU()
dropout = nn.Dropout(0.5)
layer2 = nn.Linear(256, 10)

# Manual composition
x = layer1(x)
x = activation(x)
x = dropout(x) # Automatically uses model.training state
output = layer2(x)

# Or use Sequential
model = nn.Sequential(
    nn.Linear(784, 256), nn.ReLU(), nn.Dropout(0.5),
    nn.Linear(256, 10)
)
output = model(x)

# Collect parameters
params = list(model.parameters())

```

Let's walk through each difference:

- **Line 1-2 (Import):** Both frameworks provide layers in a dedicated module. TinyTorch uses `tinytorch.core.layers`; PyTorch uses `torch.nn`.
- **Line 4-7 (Layer Creation):** Identical API. Both use `Linear(in_features, out_features)` and `Dropout(p)`.



- **Line 9-13 (Manual Composition):** TinyTorch requires explicit `training=True` flag for Dropout; PyTorch uses global model state (`model.train()`).
- **Line 15-19 (Sequential):** Identical pattern for composing layers into a container.
- **Line 22 (Parameters):** Both use `.parameters()` method to collect all trainable tensors. PyTorch returns a generator; TinyTorch returns a list.

#### 💡 What's Identical

Layer initialization API, forward pass mechanics, and parameter collection patterns. When you debug PyTorch shape errors or parameter counts, you'll understand exactly what's happening because you built the same abstractions.

### 7.9.1 Why Layers Matter at Scale

To appreciate why layer design matters, consider the scale of modern ML systems:

- **GPT-3:** 175 billion parameters across 96 Linear layers (each layer transforming 12,288 features) = **350 GB** of parameter memory
- **ResNet-50:** 25.5 million parameters with 50 convolutional and linear layers = **100 MB** of parameter memory
- **BERT-Base:** 110 million parameters with 12 transformer blocks (each containing multiple Linear layers) = **440 MB** of parameter memory

Every Linear layer in these architectures runs the same  $y = xW + b$  you just wrote. Parameter counts, memory scaling, and initialization aren't academic concerns — they're the first three things an engineer profiles when a model won't fit, won't converge, or won't run fast enough. When GPT-3 misbehaves, the people fixing it are debugging the same abstractions you built today.

## 7.10 Check Your Understanding

#### 💡 Check Your Understanding — Layers

Before moving on, verify you can articulate each of the following:

- Why parameter count for `Linear(in, out)` is `in*out + out` and how this scales linearly when you double `out_features`.
- Why all-zero (or any constant) initialization freezes training, and how LeCun-style `sqrt(1/in_features)` scaling breaks symmetry while keeping variance stable.
- Why Dropout scales surviving activations by `1/(1-p)` during training — and what goes wrong at inference time if you forget the scale.
- Why Linear's `matmul` is compute-bound (high arithmetic intensity) while Dropout and bias-add are memory-bound, and what that implies for which operation BLAS-style tiling targets.
- The role of the `parameters()` contract: why a two-method base class (forward + parameters) is enough for optimizers and autograd to treat any layer uniformly.

If any of these feels fuzzy, revisit Core Concepts (The Linear Transformation, Weight Initialization, Parameter Management, Memory and Computational Complexity) before moving on.

The collapsible Q&A below works through each of these quantitatively.

#### Q1: Parameter Scaling

A Linear layer has `in_features=784` and `out_features=256`. How many parameters does it have? If you double `out_features` to 512, how many parameters now?

 **Answer**

**Original:**  $784 \times 256 + 256 = 200,960$  parameters

**Doubled:**  $784 \times 512 + 512 = 401,920$  parameters

Doubling `out_features` roughly doubles the parameter count because weights dominate (200,704 vs 401,408 for weights alone). Parameter count scales linearly with layer width.

**Memory:**  $200,960 \times 4 = 803,840$  bytes  $\approx 785$  KB (original) vs  $401,920 \times 4 = 1,607,680$  bytes  $\approx 1.53$  MB (doubled).

**Q2: Multi-layer Memory**

A 3-layer network has architecture  $784 \rightarrow 256 \rightarrow 128 \rightarrow 10$ . Calculate total parameter count and memory usage (assume float32).

 **Answer**

**Layer 1:**  $784 \times 256 + 256 = 200,960$  parameters **Layer 2:**  $256 \times 128 + 128 = 32,896$  parameters **Layer 3:**  $128 \times 10 + 10 = 1,290$  parameters

**Total:** 235,146 parameters

**Memory:**  $235,146 \times 4 = 940,584$  bytes  $\approx 919$  KB

That's parameter memory only. Add activation memory for batch processing: at batch size 32, you also hold intermediate tensors at each layer ( $32 \times 784, 32 \times 256, 32 \times 128, 32 \times 10 \approx 147$  KB more).

**Q3: Dropout Scaling**

Why do we scale surviving values by  $1/(1-p)$  during training? What happens if we don't scale?

 **Answer**

**With scaling:** Expected value of output matches input. If  $p=0.5$ , half the neurons survive and are scaled by 2.0, so  $E[\text{output}] = 0.5 \times 0 + 0.5 \times 2x = x$ .

**Without scaling:** Expected value is halved.  $E[\text{output}] = 0.5 \times 0 + 0.5 \times x = 0.5x$ . During inference (no dropout), output would be  $x$ , creating a mismatch.

**Result:** Network sees different magnitude activations during training vs inference, leading to poor test performance. Scaling ensures consistent magnitudes.

**Q4: Computational Bottleneck**

For a Linear layer's forward pass  $y = xW + b$ , which operation dominates: matrix multiply or bias addition?

 **Answer**

**Matrix multiply:**  $O(\text{batch} \times \text{in\_features} \times \text{out\_features})$  operations **Bias addition:**  $O(\text{batch} \times \text{out\_features})$  operations

For `Linear(784, 256)` with batch size 32:

- **Matmul:**  $32 \times 784 \times 256 = 6,422,528$  operations
- **Bias:**  $32 \times 256 = 8,192$  operations

Matrix multiply dominates by  $\sim 784\times$ . That ratio is why every serious framework leans on optimized BLAS kernels (MKL, OpenBLAS) on CPU and cuBLAS on GPU — speeding up matmul speeds up the whole network.

**Q5: Initialization Impact**

What happens if you initialize all weights to zero? To the same non-zero value?

 Answer

**All zeros:** Network can't learn. All neurons compute identical outputs, receive identical gradients, and update identically. Symmetry is never broken. Training is stuck.

**Same non-zero value (e.g., all 1s):** Same problem - symmetry. All neurons remain identical throughout training. You need randomness to break symmetry.

**Proper initialization:** Random values scaled by `sqrt(1/in_features)` break symmetry AND maintain stable gradient variance. This is why proper initialization is essential for learning.

**Q6: Batch Size vs Throughput**

From your timing analysis, batch size 32 processes 10,000 samples/sec, while batch size 1 processes 800 samples/sec. Why is batching faster?

 Answer

**Overhead amortization:** Setting up matrix operations has fixed cost per call. With batch=1, you pay this cost for every sample. With batch=32, you pay once for 32 samples.

**Vectorization:** Modern CPUs/GPUs process vectors efficiently. Matrix operations on larger matrices utilize SIMD instructions and better cache locality.

**Throughput calculation:** - Batch=1: 800 samples/sec means each forward pass takes ~1.25ms - Batch=32: 10,000 samples/sec means each forward pass takes ~3.2ms for 32 samples = 0.1ms per sample

Batching achieves 12.5x better per-sample performance by better utilizing hardware.

**Trade-off:** Larger batches increase latency (time to process one sample) but dramatically improve throughput (samples processed per second).

## 7.11 Key Takeaways

- **A layer is a function with weights, plus a parameter list:** the minimal contract — `forward(x)` and `parameters()` — is all autograd and optimizers ever need from any layer, from `Linear` to transformer blocks.
- **$y = xW + b$  is the workhorse:** Linear layers dominate both parameter count and FLOPs in fully-connected stacks; every scaling law ultimately argues about how many `Linear` multiplies fit on the hardware.
- **Initialization is not a detail:** LeCun-style `sqrt(1/fan_in)` scaling keeps activation variance roughly constant through depth. Zero or uniform init freezes symmetry and blocks learning outright.
- **Dropout's inverted scaling preserves expected magnitude:** dividing survivors by `1-p` keeps train-time and inference-time activation scales aligned — skip the scaling and your model tests far worse than it trains.
- **Sequential is twelve lines but unlocks the whole zoo:** chaining `forward()` and flattening per-layer `parameters()` is the entire mechanism behind every model architecture in this book.

**Coming next:** Module 04 adds the loss functions — `MSELoss`, `CrossEntropyLoss`, `BinaryCrossEntropyLoss` — that turn your predictions and targets into the scalar signal autograd will backpropagate into these parameters.

## 7.12 Further Reading

The design of modern neural network layers represents a delicate balance between mathematical theory and hardware reality. From breaking symmetries during initialization to regularizing massive parameter spaces,

the following papers document the critical breakthroughs that made deep learning architectures scalable and stable.

### 7.12.1 Seminal Papers

- **Understanding the difficulty of training deep feedforward neural networks** - Glorot and Bengio (2010). Introduces Xavier/Glorot initialization and analyzes why proper weight scaling matters for gradient flow. The foundation for modern initialization schemes. [PMLR](#)
- **Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification** - He et al. (2015). Introduces He initialization tailored for ReLU activations. Shows how initialization schemes must match activation functions for optimal training. **Systems Implication:** Proper initialization prevented vanishing gradients in very deep networks like ResNet, allowing models to scale to hundreds of layers and fully utilize the compute capacity of modern accelerators. [arXiv:1502.01852](#)
- **Dropout: A Simple Way to Prevent Neural Networks from Overfitting** - Srivastava et al. (2014). The original dropout paper demonstrating how random neuron dropping prevents overfitting. Includes theoretical analysis and extensive empirical validation. [JMLR](#)

### 7.12.2 Additional Resources

- **Textbook:** “[Deep Learning](#)” by Goodfellow, Bengio, and Courville - Chapter 6 covers feedforward networks and linear layers in detail
- **Documentation:** [PyTorch nn.Linear](#) - See how production frameworks implement the same concepts
- **Blog Post:** “[A Recipe for Training Neural Networks](#)” by Andrej Karpathy - Practical advice on initialization, architecture design, and debugging

## 7.13 What’s Next

### Coming Up: Module 04 — Losses

Your layers can produce predictions, but you have no way to say *how wrong* a prediction is. Module 04 introduces loss functions — `MSELoss` for regression, `CrossEntropyLoss` for classification — that turn a prediction and a target into a single scalar. That scalar becomes the signal autograd (Module 06) backpropagates, which optimizers (Module 07) then apply to the very `parameters()` you collected here.

#### Preview — how your layers plug into the next modules:

Table 7.8 traces how this module is reused by later parts of the curriculum.

Table 7.8: How the Layer abstractions feed into later training modules.

| Module                | What It Does                     | Your Layers In Action                                                  |
|-----------------------|----------------------------------|------------------------------------------------------------------------|
| <b>04: Losses</b>     | Quantify prediction error        | <code>loss = CrossEntropyLoss() (model(x), y)</code>                   |
| <b>06: Autograd</b>   | Backpropagate through the stack  | <code>loss.backward()</code> fills <code>layer.weight.grad</code>      |
| <b>07: Optimizers</b> | Update parameters from gradients | <code>optimizer.step()</code> consumes <code>model.parameters()</code> |

## 7.14 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## 🔥 Chapter 8

# Module 04: Losses

A loss function reduces a full output tensor to a single scalar. That reduction is the only signal the optimizer ever sees. Build the reduction wrong, or apply it to the wrong axis, and every gradient flowing backward through autograd optimizes the wrong objective.

### Module Info

**FOUNDATION TIER** | Difficulty: ●●○○ | Time: 4-6 hours | Prerequisites: 01, 02, 03

**Prerequisites:** Modules 01 (Tensor), 02 (Activations), and 03 (Layers) must be completed. This module assumes you understand:

- Tensor operations and broadcasting (Module 01)
- Activation functions and their role in neural networks (Module 02)
- Layers and how they transform data (Module 03)

If you can build a simple neural network that takes input and produces output, you're ready to learn how to measure its quality.

## 8.1 Overview

A neural network without a loss function is a guess machine. The loss is the single scalar that tells optimization which direction to move — turning a forward pass into a learning step. Get it wrong and training stalls, diverges, or silently optimizes the wrong objective.

In this module you'll implement three losses that cover most supervised learning: **Mean Squared Error** for regression, **CrossEntropy** for multi-class classification, and **Binary Cross-Entropy** for multi-label or binary decisions. Along the way you'll implement the log-sum-exp trick — the one numerical safeguard that separates a softmax that trains from one that returns `nan` on the first batch with large logits.

By the end, you'll know not just *how* to compute each loss, but *why* the choice of loss reshapes what your model learns, and where naive implementations break at production scale.

## 8.2 Learning Objectives

### By completing this module, you will:

- **Implement** `MSELoss` for regression, `CrossEntropyLoss` for multi-class classification, and `BinaryCrossEntropyLoss` for binary decisions
- **Master** the log-sum-exp trick for numerically stable softmax computation
- **Understand** computational complexity ( $O(B \times C)$  for cross-entropy with large vocabularies) and memory trade-offs
- **Analyze** loss function behavior across different prediction patterns and confidence levels
- **Connect** your implementation to production PyTorch patterns and engineering decisions at scale

## 8.3 What You'll Build




Figure 8.1: **TinyTorch Loss Functions:** MSE for regression, and Cross-Entropy for classification tasks.

### Implementation roadmap:

Table 8.1 lays out the implementation in order, one part at a time.

Table 8.1: **Implementation roadmap for the three loss functions.**

| Step | What You'll Implement                         | Key Concept                                            |
|------|-----------------------------------------------|--------------------------------------------------------|
| 1    | <code>log_softmax()</code>                    | Log-sum-exp trick for numerical stability              |
| 2    | <code>MSELoss.forward()</code>                | Mean squared error for continuous predictions          |
| 3    | <code>CrossEntropyLoss.forward()</code>       | Negative log-likelihood for multi-class classification |
| 4    | <code>BinaryCrossEntropyLoss.forward()</code> | Cross-entropy specialized for binary decisions         |

### The pattern you'll enable:

```
# Measuring prediction quality
loss = criterion(predictions, targets) # Scalar feedback signal for learning
```

### 8.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Gradient computation (automatic differentiation is a later module)
- Advanced loss variants (Focal Loss, Label Smoothing, Huber Loss)
- Hierarchical or sampled softmax for large vocabularies (PyTorch optimization)
- Custom reduction strategies beyond mean

**You are building the core feedback signal.** Gradient-based learning comes next.



## 8.4 API Reference

This section provides a quick reference for the loss functions you'll build. Use it as your cheat sheet while implementing and debugging.

### 8.4.1 Helper Functions

```
log_softmax(x: Tensor, dim: int = -1) -> Tensor
```

Computes numerically stable log-softmax using the log-sum-exp trick. This is the foundation for cross-entropy loss.

**Parameters:** -  $x$  (Tensor): Input tensor containing logits (raw model outputs, unbounded values) -  $dim$  (int): Dimension along which to compute log-softmax (default: -1, last dimension)

**Returns:** Tensor with same shape as input, containing log-probabilities

**Note:** Logits are raw, unbounded scores from your model before any activation function. CrossEntropyLoss expects logits, not probabilities.

### 8.4.2 Loss Functions

All loss functions follow the same pattern:

Table 8.2 gives the mathematical form, output range, and typical use case of each.

Table 8.2: Constructor, forward signature, and use case for each loss.

| Loss Function          | Constructor              | Forward Signature                                       | Use Case                   |
|------------------------|--------------------------|---------------------------------------------------------|----------------------------|
| MSELoss                | MSELoss()                | forward(predictions: Tensor, targets: Tensor) -> Tensor | Regression                 |
| CrossEntropyLoss       | CrossEntropyLoss()       | forward(logits: Tensor, targets: Tensor) -> Tensor      | Multi-class classification |
| BinaryCrossEntropyLoss | BinaryCrossEntropyLoss() | forward(predictions: Tensor, targets: Tensor) -> Tensor | Binary classification      |

#### Common Pattern:

```
loss_fn = MSELoss()
loss = loss_fn(predictions, targets) # __call__ delegates to forward()
```

### 8.4.3 Input/Output Shapes

Understanding input shapes is crucial for correct loss computation:

Table 8.3 pins down the shape contract each loss expects.

Table 8.3: Input and output shape contracts for each loss function.

| Loss               | Predictions Shape            | Targets Shape                      | Output Shape |
|--------------------|------------------------------|------------------------------------|--------------|
| MSE                | $(N, )$ or $(N, D)$          | Same as predictions                | $()$ scalar  |
| CrossEntropy       | $(N, C)$ logits <sup>1</sup> | $(N, )$ class indices <sup>2</sup> | $()$ scalar  |
| BinaryCrossEntropy | probabilities <sup>3</sup>   | $(N, )$ binary labels (0 or 1)     | $()$ scalar  |

Where  $N$  = batch size,  $D$  = feature dimension,  $C$  = number of classes

**Notes:** 1. **Logits:** Raw unbounded values from your model (e.g.,  $[2.3, -1.2, 5.1]$ ). Do NOT apply softmax before passing to CrossEntropyLoss. 2. **Class indices:** Integer values from 0 to  $C-1$  indicating the correct class (e.g.,  $[0, 2, 1]$  for 3 samples). 3. **Probabilities:** Values between 0 and 1 after applying sigmoid activation. Must be in valid probability range.

## 8.5 Core Concepts

This section covers the fundamental ideas you need to understand loss functions deeply. These concepts apply to every ML framework, not just TinyTorch.

### 8.5.1 Loss as a Feedback Signal

A loss function turns “how good is my model?” into a single number that optimization can act on. Predict \$250,000 for a house that sold for \$245,000 — how wrong is that compared to predicting \$150,000? The loss assigns a precise penalty to each, and the gradient of that penalty tells the optimizer which way to move every parameter.

That second part is what matters. A loss must be differentiable: you need not just the current error, but the direction to move to reduce it. This is why MSE squares the error rather than taking the absolute value — the square has a smooth gradient everywhere, while  $|x|$  has a kink at zero that confuses optimizers.

Every training iteration is the same loop: forward pass produces predictions, loss measures error, backward pass turns that error into parameter updates. The loss value is the scalar summary of model quality for the whole batch — and the only signal the optimizer ever sees.

### 8.5.2 Mean Squared Error

MSE is the foundational loss for regression problems. It measures the average squared distance between predictions and targets. The squaring serves three purposes: it makes all errors positive (preventing cancellation), it heavily penalizes large errors, and it creates smooth gradients for optimization.

Here’s the complete implementation from your module:

The code in `?@lst-04-losses-mse` makes this concrete.

```
def forward(self, predictions: Tensor, targets: Tensor) -> Tensor:
    """Compute mean squared error between predictions and targets."""
    # Step 1: Compute element-wise difference
    diff = predictions.data - targets.data

    # Step 2: Square the differences
    squared_diff = diff ** 2

    # Step 3: Take mean across all elements
```

```
mse = np.mean(squared_diff)

return Tensor(mse)
```

: Listing 4.1 — Mean squared error forward pass: subtract, square, average. {#lst-04-losses-mse}

Three operations: subtract, square, average. Yet squaring creates a quadratic error landscape: an error of 10 contributes 100 to the loss; an error of 20 contributes 400. The optimizer sees the worst-predicted samples as four times more important than samples half as wrong, so it spends most of its capacity fixing the largest errors first.

Back to house prices. An error of \$5,000 squares to 25M. An error of \$50,000 squares to 2.5B — one hundred times the penalty for ten times the error. That asymmetry is MSE’s defining feature: it aggressively corrects outliers, but it also lets noisy labels dominate the gradient. If your dataset has bad labels, MSE will chase them.

### 8.5.3 Cross-Entropy Loss

Cross-entropy measures *surprise*: how unexpected the true label is under the distribution your model predicted. Where MSE measures geometric distance in output space, cross-entropy measures information-theoretic mismatch in probability space — and that distinction is why classifiers train faster with cross-entropy than with MSE on the same logits.

The formula reduces to a single number: the negative log-probability the model assigned to the correct class. Easy to write, easy to break — the implementation lives or dies on numerical stability:

The code in ?@lst-04-losses-cross-entropy makes this concrete.

```
def forward(self, logits: Tensor, targets: Tensor) -> Tensor:
    """Compute cross-entropy loss between logits and target class indices."""
    # Step 1: Compute log-softmax for numerical stability
    log_probs = log_softmax(logits, dim=-1)

    # Step 2: Select log-probabilities for correct classes
    batch_size = logits.shape[0]
    target_indices = targets.data.astype(int)

    # Select correct class log-probabilities using advanced indexing
    selected_log_probs = log_probs.data[np.arange(batch_size), target_indices]

    # Step 3: Return negative mean (cross-entropy is negative log-likelihood)
    cross_entropy = -np.mean(selected_log_probs)

    return Tensor(cross_entropy)
```

: Listing 4.2 — CrossEntropyLoss forward pass fusing log-softmax with target-class selection. {#lst-04-losses-cross-entropy}

The critical detail is calling `log_softmax` directly instead of `log(softmax(...))`. The two are mathematically identical but computationally worlds apart: a logit of 100 makes  $\exp(100) \approx 2.7 \times 10^{43}$ , which overflows float32 to `inf`, which becomes `nan` after division. The fused log-softmax never materializes the dangerous exponentials — see the next subsection for the trick.

Cross-entropy’s gradient pressure is asymmetric on purpose. Predict 0.99 for the correct class and the loss is  $-\log(0.99) \approx 0.01$ ; predict 0.01 and the loss is  $-\log(0.01) \approx 4.6$  — 460× larger for the same gap

in probability. The further the model is from “confident and right,” the harder cross-entropy pulls — which is exactly the dynamic you want for classification.

**Complexity.** Cross-entropy loss is  $O(B \times C)$  in both time and memory, where  $B$  is batch size and  $C$  is the number of classes. Every logit must be exponentiated (to form softmax) and every log-probability must be read back out to pick the target’s score. MSE and BCE are  $O(B \times D)$  where  $D$  is the feature dimension — cheaper, because there is no normalizing sum over classes. This is why large-vocabulary language models ( $C = 50,000+$ ) spend a disproportionate share of training time in the loss, and why hierarchical and sampled softmax variants exist to drop effective  $C$ .

#### 8.5.4 Numerical Stability in Loss Computation

The log-sum-exp trick is the single most important numerical-stability technique in classification. The problem it solves is structural: softmax requires exponentiating every logit, but `exp` overflows float32 at inputs above ~88. Real models routinely produce logits in the hundreds.

Watch what happens without the trick. Naive softmax computes  $\exp(x) / \sum(\exp(x))$ . With logits `[100, 200, 300]`, the denominator alone needs  $\exp(300) \approx 1.97 \times 10^{130}$  — `inf` in float32, and `inf / inf = nan` everywhere downstream. The fix is to subtract the per-row max *before* exponentiating, which leaves the result mathematically unchanged but pulls every exponent into the safe range:

The code in `?lst-04-losses-log-softmax` makes this concrete.

```
def log_softmax(x: Tensor, dim: int = -1) -> Tensor:
    """Compute log-softmax with numerical stability."""
    # Step 1: Find max along dimension for numerical stability
    x_max = np.max(x.data, axis=dim, keepdims=True)

    # Step 2: Subtract max to prevent overflow
    shifted = x.data - x_max

    # Step 3: Compute log(sum(exp(shifted)))
    log_sum_exp = np.log(np.sum(np.exp(shifted), axis=dim, keepdims=True))

    # Step 4: Return log_softmax = shifted - log_sum_exp
    result = shifted - log_sum_exp

    return Tensor(result)
```

: Listing 4.3 — Numerically stable log-softmax via the log-sum-exp shift. `{#lst-04-losses-log-softmax}`

After subtracting the max (300), the shifted logits become `[-200, -100, 0]`. The largest exponent is now  $\exp(0) = 1.0$  — always safe. The smallest,  $\exp(-200)$ , underflows to 0, but those terms were going to round to nothing in the sum anyway, so the loss is exact in the regime we care about.

The trick is algebraically exact, not an approximation: subtracting a constant  $m$  from every logit multiplies both the numerator and denominator of softmax by  $\exp(-m)$ , which cancels. What changes is the floating-point dynamic range — and that change is the difference between training and `nan`.

#### 8.5.5 Reduction Strategies

All three loss functions reduce a batch of per-sample errors to a single scalar by taking the mean. Mathematically, this is trivial: just sum the values and divide by the count. However, from a systems perspective, condensing millions of concurrent loss calculations down to a single number presents a massive synchronization bottleneck. When thousands of GPU threads attempt to update a global sum simultaneously, the hardware must meticulously orchestrate the memory access to prevent threads from overwriting each other’s work.

**i Systems Implication: Parallel Reductions & Atomic Locks**

To prevent data races during mean or sum calculations, hardware could use **atomic locks** to ensure only one thread updates the global sum at a time. However, excessive locking bottlenecks performance. Instead, frameworks use **parallel reductions**: threads sum their local chunks independently, and the results are aggregated hierarchically (like a tournament bracket). Furthermore, when summing millions of elements, the running total can become so large that adding small individual losses gets swallowed by **FP32 precision limits**. To prevent precision loss, these reductions are often accumulated in higher precision (FP64) or calculated in blocks.

Mean reduction does two useful things. It normalizes by batch size, so loss values are comparable across batches of 32 and 128 samples — and it makes the per-sample gradient contribution  $1/B$ , which keeps gradient magnitudes stable as you scale the batch. Sum reduction (`np.sum`) drops the  $1/B$ , so the gradient grows linearly with batch size; if you switch from `mean` to `sum`, you must shrink the learning rate by the same factor to get equivalent updates. No reduction (`reduction='none'` in PyTorch) returns the per-sample loss vector — useful for weighted sampling, hard-example mining, or per-sample diagnostics.

The reason `mean` is the universal default: it makes learning rates transferable across batch sizes without manual rescaling.

## 8.6 Common Errors

### 8.6.1 Shape Mismatch in Cross-Entropy

**Error:** `IndexError: index 5 is out of bounds for axis 1 with size 3`

This happens when your target class indices exceed the number of classes in your logits. If you have 3 classes (indices 0, 1, 2) but your targets contain index 5, the indexing operation fails.

**Fix:** Verify your target indices match your model's output dimensions. For a 3-class problem, targets should only contain 0, 1, or 2.

```
# Wrong - target index 5 doesn't exist for 3 classes
logits = Tensor([[1.0, 2.0, 3.0]]) # 3 classes
targets = Tensor([5]) # Index out of bounds

# Correct - target indices match number of classes
logits = Tensor([[1.0, 2.0, 3.0]])
targets = Tensor([2]) # Index 2 is valid for 3 classes
```

### 8.6.2 NaN Loss from Numerical Instability

**Error:** `RuntimeWarning: invalid value encountered in log` followed by `loss.data = nan`

This occurs when probabilities reach exactly 0.0 or 1.0, causing  $\log(0) = -\infty$ . Binary cross-entropy is particularly vulnerable because it computes both  $\log(\text{prediction})$  and  $\log(1-\text{prediction})$ .

**Fix:** Clamp probabilities to a safe range using epsilon:

```
# Already implemented in your BinaryCrossEntropyLoss:
eps = 1e-7
clamped_preds = np.clip(predictions.data, eps, 1 - eps)
```

This ensures you never compute  $\log(0)$  while keeping values extremely close to the true probabilities.

### 8.6.3 Confusing Logits and Probabilities

**Error:** `loss.data = inf` or unreasonably large loss values

Cross-entropy expects raw logits (unbounded values from your model), while binary cross-entropy expects probabilities (0 to 1 range). Mixing these up causes numerical explosions.

**Fix:** Check what your model outputs:

```
# CrossEntropyLoss: Use raw logits (no sigmoid/softmax!)
logits = linear_layer(x) # Raw outputs like [2.3, -1.2, 5.1]
loss = CrossEntropyLoss()(logits, targets)

# BinaryCrossEntropyLoss: Use probabilities (apply sigmoid!)
logits = linear_layer(x)
probabilities = sigmoid(logits) # Converts to [0, 1] range
loss = BinaryCrossEntropyLoss()(probabilities, targets)
```

## 8.7 Production Context

### 8.7.1 Your Implementation vs. PyTorch

Mathematically, your TinyTorch losses and PyTorch's are the same function — same formulas, same log-sum-exp safeguard. The gap is engineering: GPU kernels, fused ops, lower-precision arithmetic, and configurable reductions. The table makes the gap concrete:

Table 8.4 places your implementation side by side with the production reference for direct comparison.

Table 8.4: Feature comparison between TinyTorch losses and PyTorch equivalents.

| Feature                    | Your Implementation | PyTorch                           |
|----------------------------|---------------------|-----------------------------------|
| <b>Backend</b>             | NumPy (Python)      | C++/CUDA                          |
| <b>Numerical Stability</b> | Log-sum-exp trick   | Same trick, fused kernels         |
| <b>Speed</b>               | 1x (baseline)       | 10-100x faster (GPU)              |
| <b>Reduction Modes</b>     | Mean only           | mean, sum, none                   |
| <b>Advanced Variants</b>   |                     | Label smoothing, weights          |
| <b>Memory Efficiency</b>   | Standard            | Fused operations<br>reduce copies |

### 8.7.2 Code Comparison

The following comparison shows equivalent loss computations in TinyTorch and PyTorch. Notice how the high-level API is nearly identical - you're learning the same patterns used in production.

## 8.8 Your TinyTorch

```
from tinytorch.core.tensor import Tensor
from tinytorch.core.losses import MSELoss, CrossEntropyLoss

# Regression
mse_loss = MSELoss()
predictions = Tensor([200.0, 250.0, 300.0])
targets = Tensor([195.0, 260.0, 290.0])
loss = mse_loss(predictions, targets)

# Classification
ce_loss = CrossEntropyLoss()
logits = Tensor([[2.0, 0.5, 0.1], [0.3, 1.8, 0.2]])
labels = Tensor([0, 1])
loss = ce_loss(logits, labels)
```

## 8.9 PyTorch

```
import torch
import torch.nn as nn

# Regression
mse_loss = nn.MSELoss()
predictions = torch.tensor([200.0, 250.0, 300.0])
targets = torch.tensor([195.0, 260.0, 290.0])
loss = mse_loss(predictions, targets)

# Classification
ce_loss = nn.CrossEntropyLoss()
logits = torch.tensor([[2.0, 0.5, 0.1], [0.3, 1.8, 0.2]])
labels = torch.tensor([0, 1])
loss = ce_loss(logits, labels)
```

Let's walk through the key similarities and differences:

- **Line 1 (Imports):** Both frameworks use modular imports. TinyTorch exposes loss functions from `core.losses`; PyTorch uses `torch.nn`.
- **Line 3 (Construction):** Both use the same pattern: instantiate the loss function once, then call it multiple times. No parameters needed for basic usage.
- **Line 4-5 (Data):** TinyTorch wraps Python lists in `Tensor`; PyTorch uses `torch.tensor()`. The data structure concept is identical.
- **Line 6 (Computation):** Both compute loss by calling the loss function object. Under the hood, this calls the `forward()` method you implemented.
- **Line 9 (Classification):** Both expect raw logits (not probabilities) for cross-entropy. The `log_softmax` computation happens internally in both frameworks.

 **What's Identical**

The mathematical formulas, numerical stability techniques (log-sum-exp trick), and high-level API patterns. When you debug PyTorch loss functions, you'll understand exactly what's happening because you built the same abstractions.

**8.9.1 Why Loss Functions Matter at Scale**

To appreciate why loss functions matter in production, consider the scale of modern ML systems:

- **Language models:** 50,000-token vocabulary  $\times$  128 batch size = **6.4M exponential operations per loss computation**. Sampled softmax cuts this to  $\sim 128K$  ( $50\times$  speedup).
- **Computer vision:** ImageNet with 1,000 classes runs **256,000 softmax computations per batch**. Fused CUDA kernels drop this from  $\sim 15$  ms to  $\sim 0.5$  ms.
- **Recommendation systems:** Billions of items demand specialized losses. YouTube's recommender uses **sampled softmax over 1M+ videos**, making loss computation the primary training bottleneck.

Memory tells the same story. A language-model forward pass might burn 8 GB on activations and 2 GB on parameters, but **another 73.2 MB just for the cross-entropy loss tensors** ( $B=128$ ,  $C=50,000$ , float32 — three copies: logits, softmax, log-softmax). FP16 halves that to 36.6 MB. Hierarchical or sampled softmax avoids materializing the full vocabulary at all.

The loss computation typically accounts for **5-10% of total training time** in well-optimized systems, but can dominate (30-50%) for large vocabularies without optimization. This is why production frameworks invest heavily in fused kernels, specialized data structures, and algorithmic improvements like hierarchical softmax.

**8.10 Check Your Understanding** **Check Your Understanding — Losses**

Before moving on, verify you can articulate each of the following:

- Why MSE squares the error (positivity + smooth gradient + outlier amplification) and when that aggressive outlier weighting becomes a liability with noisy labels.
- Why `CrossEntropyLoss` expects raw logits, not probabilities — and why calling `log(softmax(...))` materializes the same `inf/nan` that `log_softmax(...)` structurally avoids.
- How the log-sum-exp shift is algebraically exact but numerically transformative: subtracting `max(x)` pulls every exponent into a safe range without changing the answer.
- Why cross-entropy is  $O(B \times C)$  in time and memory, and why a 50k-token vocabulary forces hierarchical or sampled softmax in production.
- When to pick `BinaryCrossEntropyLoss` (independent binary decisions) vs. `CrossEntropyLoss` (mutually exclusive classes), and why mixing them up yields silently wrong probability semantics.

If any of these feels fuzzy, revisit Core Concepts (Mean Squared Error, Cross-Entropy Loss, Numerical Stability in Loss Computation) before moving on.

The collapsible Q&A below works each of these through with production-scale numbers.

**Q1: Memory Calculation - Large Vocabulary Language Model**



A language model with 50,000 token vocabulary uses CrossEntropyLoss with batch size 128. Using float32, how much memory does the loss computation require for logits, softmax probabilities, and log-probabilities?

### 💡 Answer

#### Calculation:

- Logits:  $128 \times 50,000 \times 4 \text{ bytes} = \mathbf{24.4 \text{ MB}}$
- Softmax probabilities:  $128 \times 50,000 \times 4 \text{ bytes} = \mathbf{24.4 \text{ MB}}$
- Log-softmax:  $128 \times 50,000 \times 4 \text{ bytes} = \mathbf{24.4 \text{ MB}}$

**Total: 73.2 MB** just for loss computation (before any model activations).

**Key insight:** Memory scales as  $B \times C$ . Doubling vocabulary doubles loss memory. This is why large language models can't afford to materialize the full vocabulary on every forward pass — they use sampled or hierarchical softmax instead.

**Production solution:** switch to FP16 (cuts the total to **36.6 MB**), or use hierarchical/sampled softmax (reduces effective  $C$  from 50,000 to  $\sim 1,000$ ).

### Q2: Complexity Analysis - Softmax Bottleneck

Your training profile shows: Forward pass 80ms, Loss computation 120ms, Backward pass 150ms. Your model has 1,000 output classes and batch size 64. Why is loss computation so expensive, and what's the fix?

### 💡 Answer

**Problem:** Loss taking 120 ms (34% of iteration time) is unusually high — the normal range is 5–10%.

**Root cause:** CrossEntropyLoss is  $O(B \times C)$ . With  $B=64$  and  $C=1,000$ , that's **64,000** exp/log operations per batch. Implemented naively (Python loops instead of vectorized NumPy), this dominates the iteration.

#### Diagnosis steps:

1. Profile within loss: is `log_softmax` the bottleneck? (Likely yes.)
2. Check vectorization: NumPy broadcasting or Python loops?
3. Check batch size: is  $B=64$  too small to amortize vectorization overhead?

#### Fixes:

- **Immediate:** use vectorized NumPy ops, not loops.
- **Better:** switch to PyTorch on CUDA — 10–50× speedup from GPU acceleration.
- **Advanced:** for  $C > 10,000$ , use hierarchical softmax (reduces complexity to  $O(B \times \log C)$ ).

**Reality check:** In optimized PyTorch on GPU, loss should be  $\sim 5\text{ms}$  for this size, not 120ms. Your implementation in pure Python/NumPy is expected to be slower, but vectorization is crucial.

### Q3: Numerical Stability - Why Log-Sum-Exp Matters

Your model outputs logits `[50, 100, 150]`. Without the log-sum-exp trick, what happens when you compute softmax? With the trick, what values are actually computed?

### 💡 Answer

**Without the trick (naive softmax):**

```
exp_vals = [exp(50), exp(100), exp(150)]
           = [5.2×1021, 2.7×103, 1.4×10 ] # Last value overflows to inf!
softmax = exp_vals / sum(exp_vals) # inf / inf = nan
```

**Result:** NaN loss, training fails.

**With log-sum-exp trick:**

```
max_val = 150
shifted = [50-150, 100-150, 150-150] = [-100, -50, 0]
exp_shifted = [exp(-100), exp(-50), exp(0)]
             = [3.7×10-22, 1.9×10-22, 1.0] # All ≤ 1.0, safe!
sum_exp = 1.0 (others negligible)
log_sum_exp = log(1.0) = 0
log_softmax = shifted - log_sum_exp = [-100, -50, 0]
```

**Result:** Valid log-probabilities, stable training.

**Key insight:** Subtracting max makes largest value 0, so  $\exp(0) = 1.0$  is always safe. Smaller values underflow to 0, but that's fine - they contribute negligibly anyway. This is why **you must use log-sum-exp for any softmax computation**.

#### Q4: Loss Function Selection - Classification Problem

You're building a medical diagnosis system with 5 disease categories. Should you use BinaryCrossEntropyLoss or CrossEntropyLoss? What if the categories aren't mutually exclusive (patient can have multiple diseases)?

#### 💡 Answer

**Case 1: Mutually exclusive diseases** (patient has exactly one)

- **Use:** CrossEntropyLoss
- **Model output:** Logits of shape (batch\_size, 5)
- **Why:** Categories are mutually exclusive — softmax ensures probabilities sum to 1.0

**Case 2: Multi-label classification** (patient can have multiple diseases)

- **Use:** BinaryCrossEntropyLoss
- **Model output:** Probabilities of shape (batch\_size, 5) after sigmoid
- **Why:** Each disease is an independent binary decision. Softmax would incorrectly force them to sum to 1.

**Example:**

```
# Mutually exclusive (one disease)
logits = Linear(features, 5)(x) # Shape: (B, 5)
loss = CrossEntropyLoss()(logits, targets) # targets: class index 0-4

# Multi-label (can have multiple)
logits = Linear(features, 5)(x) # Shape: (B, 5)
probs = sigmoid(logits) # Independent probabilities
targets = Tensor([[1, 0, 1, 0, 0], ...]) # Binary labels for each disease
loss = BinaryCrossEntropyLoss()(probs, targets)
```

**Critical medical consideration:** Multi-label is more realistic - patients often have comorbidities!

### Q5: Batch Size Impact - Memory and Gradients

You train with batch size 32, using 4GB GPU memory. You want to increase to batch size 128. Will memory usage be 16GB? What happens to the loss value and gradient quality?

#### 💡 Answer

**Memory usage:** Yes, approximately **16 GB** (4× increase).

- Loss computation scales linearly:  $4\times \text{batch} \rightarrow 4\times \text{memory}$ .
- Activations scale linearly:  $4\times \text{batch} \rightarrow 4\times \text{memory}$ .
- Model parameters: fixed regardless of batch size.

**Problem:** if your GPU only has 12 GB, training crashes with OOM.

**Loss value:** stays roughly the same (assuming similar data):

```
batch_32_loss = mean(losses[:32])
batch_128_loss = mean(losses[:128])
# Both are means over per-sample errors; with similar data they converge to
similar values.
```

**Gradient quality:** improves with larger batch.

- Batch 32: high variance, noisy gradient estimates.
- Batch 128: lower variance, smoother gradient, steadier convergence.
- Trade-off: more compute per step, fewer steps per epoch.

**Production solution - Gradient Accumulation:**

```
# Simulate batch_size=128 with only batch_size=32 memory:
for i in range(4): # 4 micro-batches
    loss = compute_loss(data[i*32:(i+1)*32])
    loss.backward() # Accumulate gradients
optimizer.step() # Update once with accumulated gradients (4*32 = 128
effective batch)
```

This gives you the gradient quality of batch 128 with only the memory cost of batch 32!

## 8.11 Key Takeaways

- **The loss is the only signal optimization sees:** every architectural choice above it — layers, activations, batch size — is wasted if the loss measures the wrong thing.
- **MSE and cross-entropy encode different priors:** MSE penalizes Euclidean distance and loves outliers; cross-entropy penalizes information-theoretic surprise and pulls hardest on confident-and-wrong predictions. Pick by task, not by habit.
- **Log-sum-exp is the safeguard that separates training from nan:** subtract `max(x)` before exponentiating, fuse `log` and `softmax` into one op, and never materialize `exp(100)`. Every production framework does this; you just implemented it.
- **Cross-entropy is  $O(B \times C)$  — and  $C$  is the expensive axis:** a 50k-vocab language-model loss can burn 73 MB per batch in float32. Sampled softmax, hierarchical softmax, and FP16 all exist to make that bill payable.
- **Mean reduction is the universal default:** dividing by batch size makes learning rates transferable across batch sizes. Switch to `sum` and you must rescale the LR by the same factor.

**Coming next:** Module 05 builds the `DataLoader` that feeds batches of `(predictions, targets)` pairs into this loss every step — turning a single-sample error into the mini-batch gradient signal that modern training relies on.

## 8.12 Further Reading

Loss functions are the steering wheel of machine learning; changing the loss changes what the model learns to prioritize. To understand how the field moved from simple least-squares regression to the nuanced, stability-focused loss landscapes that train today’s massive models, the following papers are essential reading.

### 8.12.1 Seminal Papers

- **Improving neural networks by preventing co-adaptation of feature detectors** - Hinton et al. (2012). Introduces dropout, but also discusses cross-entropy loss and its role in preventing overfitting. Understanding why cross-entropy works better than MSE for classification is fundamental. [arXiv:1207.0580](#)
- **Focal Loss for Dense Object Detection** - Lin et al. (2017). Addresses class imbalance by reshaping the loss curve to down-weight easy examples. Shows how loss function design directly impacts model performance on real problems. [arXiv:1708.02002](#)
- **When Does Label Smoothing Help?** - Müller et al. (2019). Analyzes why adding small noise to target labels (label smoothing) improves generalization. Demonstrates that loss function details matter beyond just basic formulation. [arXiv:1906.02629](#)

### 8.12.2 Additional Resources

- **Tutorial:** [Understanding Cross-Entropy Loss](#) - PyTorch documentation with mathematical details
- **Blog post:** [“The Softmax Function and Its Derivative”](#) - Excellent explanation of log-sum-exp trick and numerical stability
- **Textbook:** [“Deep Learning”](#) by Goodfellow, Bengio, and Courville - Chapter 5 covers loss functions and maximum likelihood

## 8.13 What's Next

### Coming Up: Module 05 — DataLoader

You now have a feedback signal. The next problem is feeding it: a single sample at a time is too slow, an entire dataset at once won't fit in memory, and unshuffled data trains a different model than shuffled data. Module 05 builds the **DataLoader** — batching, shuffling, and iteration — so the loss you just wrote can be averaged over `B` samples per step instead of one.

### Preview — How Your Loss Functions Get Used in Future Modules:

Table 8.5 traces how this module is reused by later parts of the curriculum.

Table 8.5: How losses feed into subsequent training modules.

| Module                | What It Does              | Your Loss In Action                                                                  |
|-----------------------|---------------------------|--------------------------------------------------------------------------------------|
| 05: <b>DataLoader</b> | Batching + shuffling      | Feeds <code>(predictions, targets)</code> pairs into your loss every step            |
| 06: <b>Autograd</b>   | Automatic differentiation | <code>loss.backward()</code> traces the loss back into parameter gradients           |
| 07: <b>Optimizers</b> | Parameter updates         | <code>optimizer.step()</code> consumes those gradients to shrink the loss            |
| 08: <b>Training</b>   | Complete training loop    | <code>loss = criterion(outputs, targets)</code> becomes the heartbeat of every epoch |

## 8.14 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## 🔥 Chapter 9

# Module 05: DataLoader

Training is I/O-bound before it is compute-bound. `DataLoader` is where the system learns to overlap disk reads, preprocessing, and accelerator compute so the GPU never sits idle. Your batch, shuffle, and collate logic decides whether an 8-GPU box runs at 8x or at 0.8x a single-GPU baseline.

### Module Info

**FOUNDATION TIER** | Difficulty: ●●○○ | Time: 3-5 hours | Prerequisites: 01-04

**Prerequisites:** You should be comfortable with tensors, activations, layers, and losses from Modules 01-04. This module introduces data loading infrastructure that will be used by autograd, optimizers, and training loops in the following modules.

## 9.1 Overview

A naive training loop reaches into a 50,000-image dataset, picks one sample, computes a gradient, and repeats. It works. It also wastes the GPU and gets the math wrong: gradients computed on a sorted sequence of samples are not the gradients you intended. Every framework solves this with the same abstraction — a `DataLoader` sitting between storage and computation, turning raw samples into shuffled, contiguous batches.

In this module you build that abstraction. A `Dataset` says how to find sample *i*. A `DataLoader` decides how many samples to group, in what order, and when to load them. The result is a single iterator that works identically on 1,000 tensors in RAM or 100 GB of JPEGs on disk — and that you will reuse, unchanged, in every later module that trains a model.

## 9.2 Learning Objectives

### By completing this module, you will:

- **Implement** the `Dataset` abstraction and `TensorDataset` for in-memory data storage
- **Build** a `DataLoader` with intelligent batching, shuffling, and memory-efficient iteration
- **Master** the Python iterator protocol for streaming data without loading entire datasets
- **Analyze** throughput bottlenecks and memory scaling characteristics with different batch sizes
- **Connect** your implementation to PyTorch data loading patterns used in production ML systems

## 9.3 What You'll Build

**Implementation roadmap:**

Table 9.1 lays out the implementation in order, one part at a time.




Figure 9.1: TinyTorch Data Pipeline: From raw dataset storage to training-ready batches.

Table 9.1: Implementation roadmap for the Dataset and DataLoader classes.

| Step | What You'll Implement       | Key Concept                             |
|------|-----------------------------|-----------------------------------------|
| 1    | Dataset abstract base class | Universal data access interface         |
| 2    | TensorDataset (Dataset)     | Tensor-based in-memory storage          |
| 3    | DataLoader.__init__()       | Store dataset, batch size, shuffle flag |
| 4    | DataLoader.__iter__()       | Index shuffling and batch grouping      |
| 5    | DataLoader._collate_batch() | Stack samples into batch tensors        |

**The pattern you'll enable:**

```

# Transform individual samples into training-ready batches
dataset = TensorDataset(features, labels)
loader = DataLoader(dataset, batch_size=32, shuffle=True)

for batch_features, batch_labels in loader:
    # batch_features: (32, feature_dim) - ready for model.forward()
    predictions = model(batch_features)

```

**9.3.1 What You're NOT Building (Yet)**

To keep this module focused, you will **not** implement:

- Multi-process data loading (PyTorch uses `num_workers` for parallel loading)
- Automatic dataset downloads (you'll use pre-downloaded data or write custom loaders)
- Prefetching mechanisms (loading next batch while GPU processes current batch)
- Custom collation functions for variable-length sequences (that's for NLP modules)

**You are building the batching foundation.** Parallel loading optimizations come later.



## 9.4 API Reference

This section provides a quick reference for the data loading classes you'll build. Use it while implementing to verify signatures and expected behavior.

### 9.4.1 Dataset (Abstract Base Class)

```
class Dataset(ABC):
    @abstractmethod
    def __len__(self) -> int

    @abstractmethod
    def __getitem__(self, idx: int)
```

The Dataset interface enforces two requirements on all subclasses:  
Table 9.2 lists the methods you need to implement.

Table 9.2: Required methods on the Dataset abstract base class.

| Method                        | Returns          | Description                                           |
|-------------------------------|------------------|-------------------------------------------------------|
| <code>__len__()</code>        | <code>int</code> | Total number of samples in dataset                    |
| <code>__getitem__(idx)</code> | Sample           | Retrieve sample at index <code>idx</code> (0-indexed) |

### 9.4.2 TensorDataset

```
TensorDataset(*tensors)
```

Wraps one or more tensors into a dataset where samples are tuples of aligned tensor slices.

**Constructor Arguments:**

- `*tensors`: Variable number of Tensor objects, all with same first dimension

**Behavior:**

- All tensors must have identical length in dimension 0 (sample dimension)
- Returns tuple (`tensor1[idx]`, `tensor2[idx]`, ...) for each sample

### 9.4.3 DataLoader

```
DataLoader(dataset, batch_size, shuffle=False)
```

Wraps a dataset to provide batched iteration with optional shuffling.

**Constructor Arguments:**

- `dataset`: Dataset instance to load from
- `batch_size`: Number of samples per batch
- `shuffle`: Whether to randomize sample order each iteration

**Core Methods:**

Table 9.3 lists the methods you need to implement.

Table 9.3: Core methods on the DataLoader class.

| Method                             | Returns                         | Description                                                                |
|------------------------------------|---------------------------------|----------------------------------------------------------------------------|
| <code>__len__()</code>             | <code>int</code>                | Number of batches (ceiling of samples divided by <code>batch_size</code> ) |
| <code>__iter__()</code>            | <code>Iterator</code>           | Returns generator yielding batched tensors                                 |
| <code>_collate_batch(batch)</code> | <code>Tuple[Tensor, ...]</code> | Stacks list of samples into batch tensors                                  |

#### 9.4.4 Data Augmentation Transforms

```
RandomHorizontalFlip(p=0.5)
RandomCrop(size, padding=4)
Compose(transforms)
```

Transform classes for data augmentation during training. Applied to individual samples before batching.

##### RandomHorizontalFlip:

- `p`: Probability of flipping (0.0 to 1.0)
- Flips images horizontally along width axis with given probability

##### RandomCrop:

- `size`: Target crop size (int for square, tuple for (H, W))
- `padding`: Pixels to pad on each side before cropping
- Standard augmentation for CIFAR-10: pads to  $40 \times 40$ , crops back to  $32 \times 32$

##### Compose:

- `transforms`: List of transform callables to apply sequentially
- Chains multiple transforms into a pipeline

## 9.5 Core Concepts

This section explains the fundamental ideas behind efficient data loading. Understanding these concepts is essential for building and debugging ML training pipelines.

### 9.5.1 Dataset Abstraction

The Dataset abstraction separates how data is stored from how it's accessed. This separation enables the same DataLoader code to work with data stored in files, databases, memory, or even generated on-demand.

The interface is deliberately minimal: `__len__()` returns the count and `__getitem__(idx)` retrieves a specific sample. A dataset backed by 50,000 JPEG files implements the same interface as a dataset with 50,000 tensors in RAM. The DataLoader doesn't care about implementation details.

Here's the complete abstract base class from your implementation:

```
class Dataset(ABC):
    """Abstract base class for all datasets."""

    @abstractmethod
    def __len__(self) -> int:
        """Return the total number of samples in the dataset."""
```

```

    pass

    @abstractmethod
    def __getitem__(self, idx: int):
        """Return the sample at the given index."""
    pass

```

The `@abstractmethod` decorator forces every subclass to implement both methods. Calling `Dataset()` raises `TypeError`, which is what you want — there is no useful default.

A minimal interface is also a composable one. A caching wrapper, a subset slicer, and a concatenation of two datasets all satisfy `__len__` and `__getitem__`, so they all plug into the same `DataLoader` without knowing or caring how the underlying samples are stored.

### 9.5.2 Batching Mechanics

Batching transforms individual samples into the stacked tensors that GPUs process efficiently. When you call `dataset[0]`, you might get `(features: (784,), label: scalar)` for an MNIST digit. When you call `next(iter(dataloader))`, you get `(features: (32, 784), labels: (32,))`. The `DataLoader` collected 32 individual samples and stacked them along a new batch dimension.

Here's how collation happens in your implementation:

The code in `?lst-05-dataloader-collate` makes this concrete.

```

def _collate_batch(self, batch: List[Tuple[Tensor, ...]]) -> Tuple[Tensor, ...]:
    """Collate individual samples into batch tensors."""
    if len(batch) == 0:
        return ()

    # Determine number of tensors per sample
    num_tensors = len(batch[0])

    # Group tensors by position
    batched_tensors = []
    for tensor_idx in range(num_tensors):
        # Extract all tensors at this position
        tensor_list = [sample[tensor_idx].data for sample in batch]

        # Stack into batch tensor
        batched_data = np.stack(tensor_list, axis=0)
        batched_tensors.append(Tensor(batched_data))

    return tuple(batched_tensors)

```

: Listing 5.1 — `_collate_batch` stacking per-position samples into contiguous batch tensors. `{#lst-05-dataloader-collate}`

The algorithm: for each position in the sample tuple (features, labels, etc.), collect all samples' values at that position, then stack them using `np.stack()` along axis 0. The result is a batch tensor where the first dimension is batch size.

Consider the memory transformation. Five individual samples might each be a `(784,)` tensor consuming 3 KB. After collation, you have a single `(5, 784)` tensor consuming 15 KB. The data is identical, but the layout is now batch-friendly: all 5 samples are contiguous in memory, enabling efficient vectorized operations.

### 9.5.3 Shuffling and Randomization

Shuffling prevents the model from learning the order of training data rather than actual patterns. Without shuffling, a model sees identical batch combinations every epoch, creating correlations between gradient updates.

The naive implementation would load all samples, shuffle the data array, then iterate. But this requires memory proportional to dataset size. Your implementation is smarter: it shuffles indices, not data.

Here's the shuffling logic from your `__iter__` method:

The code in `?@lst-05-dataloader-iter` makes this concrete.

```
def __iter__(self) -> Iterator:
    """Return iterator over batches."""
    # Create list of indices
    indices = list(range(len(self.dataset)))

    # Shuffle if requested
    if self.shuffle:
        random.shuffle(indices)

    # Yield batches
    for i in range(0, len(indices), self.batch_size):
        batch_indices = indices[i:i + self.batch_size]
        batch = [self.dataset[idx] for idx in batch_indices]

        # Collate batch
        yield self._collate_batch(batch)
```

: Listing 5.2 — `Dataloader __iter__` permuting indices and yielding batches lazily via a generator. `{#lst-05-dataloader-iter}`

The key insight: `random.shuffle(indices)` permutes a list of integers, not the underlying data. For 50,000 samples, that's 400 KB of indices instead of potentially gigabytes of images. The samples never move; only the access order changes.

A fresh shuffle each epoch means sample 42 and sample 1337 land in the same batch in epoch 1 but different batches in epoch 2. That decorrelation is what makes mini-batch SGD an unbiased estimator of the full-dataset gradient — without it, the model can fit the *order* of the data instead of the data itself.

The total cost is  $8 \text{ bytes} \times \text{dataset\_size}$ . One million samples costs 8 MB to shuffle, paid once per epoch. The shuffle itself is  $O(n)$  time, also paid once per epoch — never per batch.

### 9.5.4 Iterator Protocol and Generator Pattern

Python's iterator protocol enables `for batch in dataloader` syntax. When Python encounters this loop, it first calls `dataloader.__iter__()` to get an iterator object. Your `__iter__` method is a generator function (contains `yield`), so Python automatically creates a generator that produces values lazily.

Here's the complete implementation showing the generator pattern:

```
def __iter__(self) -> Iterator:
    """Return iterator over batches."""
    # Create list of indices
    indices = list(range(len(self.dataset)))
```

```

# Shuffle if requested
if self.shuffle:
    random.shuffle(indices)

# Yield batches - this is a generator function
for i in range(0, len(indices), self.batch_size):
    batch_indices = indices[i:i + self.batch_size]
    batch = [self.dataset[idx] for idx in batch_indices]

    # Collate batch
    yield self._collate_batch(batch)

```

Each `next()` call resumes the generator, runs until the next `yield`, hands back a batch, and pauses there. The function’s local state — `indices`, `i`, the loop variable — survives between yields, so no batch state leaks back into the caller and no caller state leaks forward into the next batch.

That laziness is the whole memory argument. At any instant, only the current batch is alive: the previous one is unreachable, the next one does not yet exist. Iterating 1,000 batches of 32 images costs the memory of 32 images, not 32,000 — a  $1,000\times$  reduction for a one-line change in API design.

The same protocol also makes infinite datasets free. A synthetic-data `__getitem__` can return a sample for any integer index, and the generator will happily yield batches forever. The *training loop* decides when to stop; the dataset never has to know.

### 9.5.5 Memory-Efficient Loading

The combination of Dataset abstraction and DataLoader iteration creates a memory-efficient pipeline regardless of dataset size.

For in-memory datasets like `TensorDataset`, all data is preloaded, but `DataLoader` still provides memory benefits by controlling how much data is active at once. Your training loop processes one batch, computes gradients, updates weights, then discards that batch before loading the next. Peak memory is `batch_size × sample_size`, not `dataset_size × sample_size`.

For disk-backed datasets, the benefits are dramatic. Consider an `ImageDataset` that loads JPEGs on-demand:

```

class ImageDataset(Dataset):
    def __init__(self, image_paths, labels):
        self.image_paths = image_paths # Just file paths (tiny memory)
        self.labels = labels

    def __len__(self):
        return len(self.image_paths)

    def __getitem__(self, idx):
        # Load image only when requested
        image = load_jpeg(self.image_paths[idx])
        return Tensor(image), Tensor(self.labels[idx])

```

When `DataLoader` calls `dataset[idx]`, the image is loaded *at that moment*, not at construction. Once the batch is consumed, the references go out of scope and the memory is reclaimed. That is why a 100 GB dataset can train on an 8 GB machine — only one batch lives in memory at a time.

This is the payoff for separating length from access. The dataset can announce that it holds 50,000 images without holding them; the DataLoader pulls exactly the indices it needs for the current batch and nothing more. Storage size and working-set size become independent.

However, this elegant on-demand loading creates a catastrophic performance cliff. Fetching data sequentially exactly when the GPU requests it introduces severe latency. The data must travel a long, slow path: from the hard disk to system RAM, then across the PCIe bus, and finally into the GPU's VRAM. While this journey happens, the GPU—capable of trillions of operations per second—sits entirely idle, starved of data.

### **i** Systems Implication: The I/O Bottleneck & PyTorch Solutions

To prevent the GPU from sitting idle while data makes the long journey from **Disk** → **RAM** → **PCIe** → **VRAM**, modern DataLoaders employ **asynchrony** and **memory pinning**. In PyTorch, practitioners solve this hardware bottleneck using two key parameters:

1. **num\_workers**: Uses multiple background *processes* to pre-fetch and decode the next batch while the GPU computes the current one. PyTorch uses processes rather than threads because the **Python GIL** (Global Interpreter Lock) prevents multiple threads from executing Python code in parallel.
2. **pin\_memory=True**: Allocates data in page-locked (pinned) system RAM. This enables the PCIe controller to perform direct memory access (DMA) transfers to the GPU VRAM much faster, without CPU intervention.

Together, multi-processing bypasses the GIL and pinning accelerates PCIe transfers, keeping the data pipeline flowing fast enough to feed the beastly GPU.

## 9.6 Common Errors

These are the most frequent mistakes encountered when implementing and using data loaders.

### 9.6.1 Mismatched Tensor Dimensions

**Error:** ValueError: All tensors must have same size in first dimension

This happens when you try to create a TensorDataset with tensors that have different numbers of samples:

```
features = Tensor(np.random.randn(100, 10)) # 100 samples
labels = Tensor(np.random.randn(90)) # 90 labels - MISMATCH!
dataset = TensorDataset(features, labels) # Raises ValueError
```

The first dimension is the sample dimension. If features has 100 samples but labels has 90, TensorDataset cannot pair them correctly.

**Fix:** Ensure all tensors have identical first dimension before constructing TensorDataset.

### 9.6.2 Forgetting to Shuffle Training Data

**Symptom:** Model converges slowly or gets stuck at suboptimal accuracy

Without shuffling, the model sees identical batch combinations every epoch. If your dataset is sorted by class (all cats, then all dogs), early batches are all cats and later batches are all dogs. The model oscillates between cat features and dog features rather than learning a unified representation.

```
# Wrong - no shuffling means same batches every epoch
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=False)
```

```
# Correct - shuffle for training
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
# But don't shuffle validation - you want consistent evaluation
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False)
```

**Fix:** Always shuffle training data, never shuffle validation or test data.

### 9.6.3 Assuming Fixed Batch Size

**Symptom:** Index errors or shape mismatches on last batch

If your dataset has 100 samples and `batch_size=32`, you get batches of size `[32, 32, 32, 4]`. The last batch is smaller because 100 is not divisible by 32. Code that assumes every batch has exactly 32 samples will fail on the last batch.

```
def train_step(batch):
    features, labels = batch
    # Wrong - assumes batch_size=32
    assert features.shape[0] == 32 # Fails on last batch!

    # Correct - get actual batch size
    batch_size = features.shape[0]
```

**Fix:** Always derive batch size from tensor shape, never hardcode it.

### 9.6.4 Index Out of Bounds

**Error:** `IndexError: Index 100 out of range for dataset of size 100`

This happens when trying to access an index that doesn't exist. Remember that Python uses 0-indexing: valid indices for a dataset of size 100 are 0 through 99, not 1 through 100.

**Fix:** Ensure index range is `0 <= idx < len(dataset)`.

## 9.7 Production Context

### 9.7.1 Your Implementation vs. PyTorch

Your `DataLoader` and PyTorch's `torch.utils.data.DataLoader` share the same conceptual design and interface. The differences are in advanced features and performance optimizations.

Table 9.4 places your implementation side by side with the production reference for direct comparison.

Table 9.4: Feature comparison between TinyTorch `DataLoader` and PyTorch `DataLoader`.

| Feature          | Your Implementation               | PyTorch                                |
|------------------|-----------------------------------|----------------------------------------|
| <b>Interface</b> | Dataset + <code>DataLoader</code> | Identical pattern                      |
| <b>Batching</b>  | Sequential in main process        | Parallel with <code>num_workers</code> |
| <b>Shuffling</b> | Index-based, $O(n)$               | Same algorithm                         |
| <b>Collation</b> | <code>np.stack()</code> in Python | Custom collate functions supported     |

| Feature            | Your Implementation | PyTorch                          |
|--------------------|---------------------|----------------------------------|
| <b>Prefetching</b> | None                | Loads next batch during compute  |
| <b>Memory</b>      | One batch at a time | Configurable buffer with workers |

### 9.7.2 Code Comparison

The following comparison shows identical usage patterns between TinyTorch and PyTorch. Notice how the APIs mirror each other exactly.

## 9.8 Your TinyTorch

```
from tinytorch.core.dataloader import TensorDataset, DataLoader

# Create dataset
features = Tensor(X_train)
labels = Tensor(y_train)
dataset = TensorDataset(features, labels)

# Create loader
train_loader = DataLoader(
    dataset,
    batch_size=32,
    shuffle=True
)

# Training loop
for epoch in range(num_epochs):
    for batch_features, batch_labels in train_loader:
        predictions = model(batch_features)
        loss = loss_fn(predictions, batch_labels)
        loss.backward()
        optimizer.step()
```

## 9.9 PyTorch

```
from torch.utils.data import TensorDataset, DataLoader

# Create dataset
features = torch.tensor(X_train)
labels = torch.tensor(y_train)
dataset = TensorDataset(features, labels)

# Create loader
```



```

train_loader = DataLoader(
    dataset,
    batch_size=32,
    shuffle=True,
    num_workers=4 # Parallel loading
)

# Training loop
for epoch in range(num_epochs):
    for batch_features, batch_labels in train_loader:
        predictions = model(batch_features)
        loss = loss_fn(predictions, batch_labels)
        loss.backward()
        optimizer.step()

```

There is exactly one substantive difference: PyTorch accepts `num_workers=4`, which spawns four worker processes that load batches in parallel and hand them back through a queue. Dataset construction and the training loop are byte-for-byte identical because they have to be — the iterator protocol is the contract, not a particular implementation of it.

#### 💡 What's Identical

The Dataset abstraction, DataLoader interface, and batching semantics are identical. When you understand TinyTorch's data pipeline, you understand PyTorch's data pipeline. The only difference is PyTorch adds parallel loading to hide I/O latency.

### 9.9.1 Why DataLoaders Matter at Scale

To see why this infrastructure earns its keep, look at production-scale training:

- **ImageNet training:** 1.2M images at  $224 \times 224 \times 3 = 600$  GB uncompressed
- **Batch memory:** `batch_size=256`  $\times$  150 KB per JPEG  $\approx$  **38 MB per batch**
- **I/O throughput:** Reading 38 MB from SSD at 500 MB/s  $\approx$  **76 ms per batch** of disk I/O alone

A forward+backward pass on the same batch takes about 50 ms on a modern GPU. Without overlap, the GPU is idle more than half of every step waiting for bytes that have not arrived yet.

Production solutions:

- **Prefetching:** Load batch N+1 while GPU processes batch N (PyTorch's `num_workers`)
- **Data caching:** Keep decoded images in RAM across epochs (eliminates JPEG decode overhead)
- **Faster formats:** Use LMDB or TFRecords instead of individual files (reduces filesystem overhead)

Your DataLoader provides the interface that enables these optimizations. Add `num_workers`, swap `TensorDataset` for a disk-backed dataset, and the training loop code stays identical.

## 9.10 Check Your Understanding

### 💡 Check Your Understanding — DataLoader

Before moving on, verify you can articulate each of the following:

- Why the `Dataset` contract is just `__len__` and `__getitem__` — and how that minimal interface lets the same `DataLoader` handle tensors in RAM or JPEGs on disk.
- Why shuffling permutes *indices*, not samples: 8 MB of integers instead of the full dataset, paid once per epoch, with no movement of the underlying bytes.
- How the generator-based `__iter__` keeps peak memory at `batch_size × sample_size` instead of `dataset_size × sample_size`, regardless of how many batches you iterate.
- Why disk-backed datasets introduce the `Disk → RAM → PCIe → VRAM` latency cliff, and how `num_workers` (process-based prefetch) and `pin_memory` (DMA-friendly allocation) hide it.
- Why `np.stack` in `_collate_batch` produces a contiguous batch tensor — and why that contiguity is what makes downstream `matmul` cache-friendly.

If any of these feels fuzzy, revisit Core Concepts (Dataset Abstraction, Shuffling and Randomization, Iterator Protocol and Generator Pattern, Memory-Efficient Loading) before moving on.

The collapsible Q&A below grounds each point in concrete ImageNet/CIFAR-scale numbers.

#### Q1: Memory Calculation

You're training on CIFAR-10 with 50,000 RGB images ( $32 \times 32 \times 3$  pixels, float32). What's the memory usage for `batch_size=128`?

#### 💡 Answer

Each image:  $32 \times 32 \times 3 \times 4 \text{ bytes} = 12,288 \text{ bytes} \approx 12 \text{ KB}$ .

Batch of 128 images:  $128 \times 12 \text{ KB} = \mathbf{1,536 \text{ KB} \approx 1.5 \text{ MB}}$ .

That is the floor — input bytes only. Add activations, gradients, and parameters and peak memory typically lands 50–100× higher. The takeaway: **batch size sets the baseline, and everything else scales from it.**

#### Q2: Throughput Analysis

Your training reports these timings per batch:

- Data loading: 45 ms
- Forward pass: 30 ms
- Backward pass: 35 ms
- Optimizer step: 10 ms

Total: 120 ms per batch. Where's the bottleneck? How much faster could training be if you eliminated data loading overhead?

#### 💡 Answer

Data loading takes 45 ms out of 120 ms = **37.5% of total time**.

If data loading were free (perfect prefetching, hot cache), total time drops to  $30+35+10 = \mathbf{75 \text{ ms per batch}}$ .

Speedup:  $120 \text{ ms} \rightarrow 75 \text{ ms} = \mathbf{1.6 \times \text{faster training}}$  from a single fix.

This is why production systems prefetch with `num_workers`: the CPU loads batch N+1 while the GPU computes batch N, and the I/O wall-time disappears under the compute.

**Q3: Shuffle Memory Overhead**

You're training on a dataset with 10 million samples. How much extra memory does `shuffle=True` require compared to `shuffle=False`?

**💡 Answer**

Index array:  $10,000,000 \times 8 \text{ bytes} = 80 \text{ MB}$ .

That is the entire overhead. The samples themselves never move.

If each sample is 10 KB, the dataset is 100 GB. Shuffling 100 GB of data costs 80 MB of indices — **0.08% overhead**. This is why every production loader shuffles indices, never bytes.

**Q4: Batch Size Trade-offs**

You're deciding between `batch_size=32` and `batch_size=256` for ImageNet training:

- `batch_size=32`: 14 hours training, 76.1% accuracy
- `batch_size=256`: 6 hours training, 75.8% accuracy

Which would you choose for a research experiment where accuracy is critical? Which for a production job where you train 100 models per day?

**💡 Answer**

**Research (accuracy critical):** `batch_size=32`

- 14 hours is acceptable for research (run overnight)
- 76.1% vs 75.8% = 0.3% accuracy gain might be significant for publication
- Smaller batches often generalize better (noisier gradients act as regularization)

**Production (throughput critical):** `batch_size=256`

- 6 hours vs 14 hours = **2.3× faster**, enabling 100 models to train in reasonable time
- 0.3% accuracy difference is negligible for many production applications
- Can try learning rate adjustments to recover accuracy while keeping speed

**Systems insight:** Batch size creates a three-way trade-off between training speed, memory usage, and model quality. The “right” answer depends on your bottleneck: time, memory, or accuracy.

**Q5: Collation Cost**

Your `DataLoader` collates batches using `np.stack()`. For `batch_size=128` with samples of shape `(3, 224, 224)`, how much data is copied during collation?

**💡 Answer**

Each sample:  $3 \times 224 \times 224 \times 4 \text{ bytes} = 602,112 \text{ bytes} \approx 588 \text{ KB}$ .

Batch of 128 samples:  $128 \times 588 \text{ KB} = 75,264 \text{ KB} \approx 73.5 \text{ MB}$ .

`np.stack()` allocates a new contiguous buffer of that size and copies all 128 samples into it. At a memory bandwidth of 20 GB/s, the copy takes **~3.7 milliseconds**.

Larger batches pay a higher *absolute* collation cost (more bytes to move) but a lower *per-sample* cost — one big copy beats 128 small ones because the memory subsystem is happiest moving long, contiguous runs.

## 9.11 Key Takeaways

- **A two-method contract scales to every dataset you will ever meet:** `__len__` + `__getitem__` is enough for the DataLoader to treat in-memory tensors, disk-backed JPEGs, or streaming sources interchangeably.
- **Shuffle indices, never bytes:** permuting an 8-byte integer array is  $O(n)$  and essentially free; permuting the data array costs dataset-sized I/O and defeats the whole point of batching.
- **Generators keep peak memory at one batch:** `yield` pauses the iterator between batches, so only the current batch is alive — which is how a 100 GB dataset trains on an 8 GB machine.
- **Collation exists to make batches contiguous:** `np.stack` produces a single flat buffer the GPU can stream, turning 128 scattered samples into one cache-friendly tensor.
- **The I/O wall is the real enemy at scale:** without prefetching, the GPU stalls at 50% utilization waiting `forDisk` → `RAM` → `PCIe` → `VRAM`; `num_workers` + `pin_memory` are what make that latency disappear.

**Coming next:** Module 06 builds automatic differentiation on top of this iterator — every batch tensor the DataLoader yields becomes a leaf of a computation graph that `loss.backward()` traces back to fill `param.grad` for every weight.

## 9.12 Further Reading

As models grew from millions to billions of parameters, the bottleneck in training shifted from simply computing gradients to keeping the monstrous hardware fed with data. The following papers illustrate how the ML community realized that data infrastructure, augmentation, and batching strategies are just as critical to model convergence as the architecture itself.

### 9.12.1 Seminal Papers

- **ImageNet Classification with Deep Convolutional Neural Networks** - Krizhevsky et al. (2012). The AlexNet paper that popularized large-scale image training and highlighted data augmentation as essential for generalization. **Systems Implication:** The model had to be split across two GTX 580 GPUs because of the strict 3GB VRAM limit, while data augmentation pipelines introduced severe CPU bottlenecks during data loading. [NeurIPS](#)
- **Accurate, Large Minibatch SGD** - Goyal et al. (2017). Facebook AI Research paper exploring how to scale batch size to 8192 while maintaining accuracy, revealing the relationship between batch size, learning rate, and convergence. [arXiv:1706.02677](#)
- **Mixed Precision Training** - Micikevicius et al. (2018). NVIDIA paper showing how batch size interacts with numerical precision for memory and speed trade-offs. [arXiv:1710.03740](#)

### 9.12.2 Additional Resources

- **Engineering Blog:** “PyTorch DataLoader Internals” — Detailed explanation of multi-process loading and prefetching strategies
- **Documentation:** [PyTorch Data Loading Tutorial](#) - See how production frameworks extend the patterns you’ve built

## 9.13 What's Next

### Coming Up: Module 06 - Autograd

You can now move data through a model. You cannot yet learn from it — every batch leaves the network as a loss number with no gradient attached. Module 06 fixes that by building automatic differentiation: every tensor remembers the operations that produced it, and `loss.backward()` walks the resulting graph to assign a gradient to every parameter the loader's batch touched.

`DataLoader` and `autograd` compose directly: the iterator you just built becomes the input edge of every computation graph in the rest of the book.

#### Where this `DataLoader` shows up next:

Table 9.5 traces how this module is reused by later parts of the curriculum.

Table 9.5: How the `DataLoader` feeds into subsequent training modules.

| Module                  | What it does                 | Your <code>DataLoader</code> in action                               |
|-------------------------|------------------------------|----------------------------------------------------------------------|
| <b>06: Autograd</b>     | Reverse-mode differentiation | Each batch tensor becomes a leaf of the computation graph            |
| <b>08: Training</b>     | End-to-end training loops    | <code>for batch in loader:</code> is the outer loop of every example |
| <b>09: Convolutions</b> | Convolutional layers         | The same iterator now feeds 4-D image batches to CNNs                |

## 9.14 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## 🔥 Chapter 10

# Module 06: Autograd

Gradient computation is the single largest source of memory growth during training. Autograd's choices about what to cache during the forward pass, what to recompute on the backward, and when to release tensors decide whether a model fits in VRAM. The runtime you build here is the memory manager for the rest of the system.

### Module Info

**FOUNDATION TIER** | Difficulty: ●●●○ | Time: 6-8 hours | Prerequisites: 01-05

You need to be fluent with everything from Modules 01–05:

- Tensor operations (matmul, broadcasting, reductions)
- Activation functions (the source of non-linearity)
- Neural network layers (what gradients will flow through)
- Loss functions (the scalar gradients flow back from)
- DataLoader for batched iteration

If you can hand-compute a forward pass through a small network and explain why we minimize loss, you're ready.

## 10.1 Overview

A neural network learns by nudging every parameter in the direction that lowers the loss. To find that direction you need a gradient — one number per parameter. A modern model has billions of parameters, so deriving those gradients by hand is not just tedious, it is impossible. Every framework you have ever used — PyTorch, TensorFlow, JAX — solves this with the same trick: automatic differentiation.

In this module you build reverse-mode autograd from scratch. The forward pass records each operation into a small graph; `loss.backward()` walks that graph in reverse, applying the chain rule one operation at a time. When you finish, calling `loss.backward()` on your tensors does the same thing it does in PyTorch — and you will know exactly why.

This is the conceptually hardest module in the Foundation tier. It is also the one that unlocks everything that follows: optimizers, training loops, and any model that learns from data.

## 10.2 Learning Objectives

 By completing this module, you will:

- **Implement** the Function base class that enables gradient computation for all operations
- **Build** computation graphs that track dependencies between tensors during forward pass
- **Master** the chain rule by implementing backward passes for arithmetic, matrix multiplication, and reductions

- **Understand** memory trade-offs between storing intermediate values and recomputing forward passes
- **Connect** your autograd implementation to PyTorch’s design patterns and production optimizations

### 10.3 What You’ll Build




Figure 10.1: **TinyTorch Autograd Engine:** Reverse-mode automatic differentiation infrastructure.

#### Implementation roadmap:

Table 10.1 lays out the implementation in order, one part at a time.

Table 10.1: **Implementation roadmap for the reverse-mode autograd engine.**

| Part | What You’ll Implement                    | Key Concept                                      |
|------|------------------------------------------|--------------------------------------------------|
| 1    | Function base class                      | Storing inputs for backward pass                 |
| 2    | AddBackward, MulBackward, MatmulBackward | Operation-specific gradient rules                |
| 3    | backward() method on Tensor              | Reverse-mode differentiation                     |
| 4    | enable_autograd() enhancement            | Monkey-patching operations for gradient tracking |
| 5    | Integration tests                        | Multi-layer gradient flow                        |

#### The pattern you’ll enable:

```

# Automatic gradient computation
x = Tensor([2.0], requires_grad=True)
y = x * 3 + 1 # y = 3x + 1
y.backward() # Computes dy/dx = 3 automatically
print(x.grad) # [3.0]
  
```

#### 10.3.1 What You’re NOT Building (Yet)

To keep this module focused, you will **not** implement:



- Higher-order derivatives (gradients of gradients)—PyTorch supports this with `create_graph=True`
- Dynamic computation graphs—your graphs are built during forward pass only
- GPU kernel fusion—PyTorch’s JIT compiler optimizes backward pass operations
- Checkpointing for memory efficiency—that’s an advanced optimization technique

**You are building the core gradient engine.** Advanced optimizations come in production frameworks.

## 10.4 API Reference

This section documents the autograd components you’ll build. These integrate with the existing `Tensor` class from Module 01.

### 10.4.1 Function Base Class

```
Function(*tensors)
```

Base class for all differentiable operations. Every operation (addition, multiplication, etc.) inherits from `Function` and implements gradient computation rules.

### 10.4.2 Core Function Classes

Table 10.2 lists the backward `Function` classes and the gradient rule each one applies.

Table 10.2: Backward Function classes and their gradient rules.

| Class                          | Purpose                         | Gradient Rule                                                                             |
|--------------------------------|---------------------------------|-------------------------------------------------------------------------------------------|
| <code>AddBackward</code>       | Addition gradients              | $\partial(a+b)/\partial a = 1, \partial(a+b)/\partial b = 1$                              |
| <code>SubBackward</code>       | Subtraction gradients           | $\partial(a-b)/\partial a = 1, \partial(a-b)/\partial b = -1$                             |
| <code>MulBackward</code>       | Multiplication gradients        | $\partial(ab)/\partial a = b, \partial(ab)/\partial b = a$                                |
| <code>DivBackward</code>       | Division gradients              | $\partial(a/b)/\partial a = 1/b, \partial(a/b)/\partial b = -a/b^2$                       |
| <code>MatmulBackward</code>    | Matrix multiplication gradients | $\partial(A@B)/\partial A = \text{grad}@B.T, \partial(A@B)/\partial B = A.T@grad$         |
| <code>SumBackward</code>       | Reduction gradients             | $\partial\text{sum}(a)/\partial a_i = 1$ for all $i$                                      |
| <code>ReshapeBackward</code>   | Shape manipulation              | $\partial(X.\text{reshape}(...))/\partial X = \text{grad}.\text{reshape}(X.\text{shape})$ |
| <code>TransposeBackward</code> | Transpose gradients             | $\partial(X.T)/\partial X = \text{grad}.T$                                                |

**Additional Backward Classes:** The implementation includes backward functions for activations (`ReLUBackward`, `SigmoidBackward`, `SoftmaxBackward`, `GELUBackward`), losses (`MSEBackward`, `BCEBackward`, `CrossEntropyBackward`), and other operations (`PermuteBackward`, `SliceBackward`). These follow the same pattern as the core classes above.

### 10.4.3 Enhanced Tensor Methods

Your implementation adds these methods to the `Tensor` class:

Table 10.3 lists the new methods autograd adds to the `Tensor` class.

Table 10.3: Methods added to the Tensor class for autograd.

| Method    | Signature                                       | Description                           |
|-----------|-------------------------------------------------|---------------------------------------|
| backward  | <code>backward(gradient=None) -&gt; None</code> | Compute gradients via backpropagation |
| zero_grad | <code>zero_grad() -&gt; None</code>             | Reset gradients to None               |

#### 10.4.4 Global Activation

Table 10.4 lists the global helpers that toggle gradient tracking.

Table 10.4: Global helper functions for enabling autograd.

| Function                     | Signature                                            | Description                         |
|------------------------------|------------------------------------------------------|-------------------------------------|
| <code>enable_autograd</code> | <code>enable_autograd(quiet=False) -&gt; None</code> | Activate gradient tracking globally |

## 10.5 Core Concepts

This section covers the fundamental ideas behind automatic differentiation. Understanding these concepts deeply will help you debug gradient issues in any framework, not just TinyTorch.

### 10.5.1 Computation Graphs

A computation graph is a directed acyclic graph (DAG): nodes are tensors, edges are the operations that produced them. When you write  $y = x * 3 + 1$ , you build a graph with three tensor nodes ( $x$ ,  $temp$ ,  $y$ ) and two operation edges (multiply, add). You don't see this graph because autograd builds it for you, silently, as a side effect of running the forward pass.

The construction trick is small but powerful: every tensor produced by an operation stores a reference to the operation that produced it. That reference — `_grad_fn` in your implementation, `grad_fn` in PyTorch — is the entire graph. To traverse the graph backward you just follow `_grad_fn` pointers until you reach the leaves.

Forward Pass:  $x \rightarrow [Mul(*3)] \rightarrow temp \rightarrow [Add(+1)] \rightarrow y$

Backward Pass:  $grad_x \leftarrow [MulBackward] \leftarrow grad_{temp} \leftarrow [AddBackward] \leftarrow grad_y$

Each backward node also has to remember the *values* it will need later. For  $z = a * b$ , the gradient with respect to  $a$  is  $grad_z * b$  — so the multiply operation must hold on to  $b$  from the forward pass. This is the central memory trade-off of autograd: every saved tensor is bytes you cannot reclaim until backward runs, but those saved tensors are exactly what makes the backward pass cheap.

Your implementation tracks graphs with the `_grad_fn` attribute:

The code in `?@lst-06-autograd-add-backward` makes this concrete.

```
class AddBackward(Function):
    """Gradient computation for addition."""

    def __init__(self, a, b):
        """Store inputs needed for backward pass."""
        self.saved_tensors = (a, b)
```

```
def apply(self, grad_output):
    """Compute gradients for both inputs."""
    return grad_output, grad_output # Addition distributes gradients equally
```

: Listing 6.1 — **AddBackward** stores inputs and distributes the incoming gradient equally to both addends. {#lst-06-autograd-add-backward}

When you compute  $z = x + y$ , your enhanced Tensor class automatically creates an `AddBackward` instance and attaches it to  $z$ :

```
result = x.data + y.data
result_tensor = Tensor(result)
result_tensor._grad_fn = AddBackward(x, y) # Track operation
```

This simple pattern scales elegantly, enabling arbitrarily complex computation graphs. However, this flexibility masks a profound structural cost: the mere act of recording operations fundamentally alters the memory lifecycle of the tensors involved.

### **i** Systems Implication: Activation Pinning

By passing  $(x, y)$  into `AddBackward(x, y)`, the computation graph captures a hard reference to the input tensors. To the Python interpreter, this means the Garbage Collector cannot free the memory for  $x$  and  $y$  until the backward pass is complete and the graph is destroyed. This “Activation Pinning” is the root cause of CUDA Out of Memory (OOM) errors during training, as every layer’s output is kept alive in VRAM for the entirety of the forward pass. In practice, researchers rely on telemetry tools like `torch.cuda.memory_allocated()` to hunt down subtle computation graph memory leaks and measure the exact overhead of retaining these activations.

Because these pinned activations iteratively consume precious VRAM, scaling networks to hundreds of layers requires rigorous strategies—such as gradient checkpointing—to mitigate this memory exhaustion while preserving the integrity of the gradient flow.

## 10.5.2 The Chain Rule

Backpropagation is the chain rule, applied one node at a time. For a composite function  $z = f(g(x))$ , the chain rule says  $dz/dx = (dz/dg) * (dg/dx)$ . Reverse-mode autograd flips this on its head: instead of multiplying derivatives left-to-right, you walk the graph backward and multiply right-to-left, so every intermediate gradient is computed exactly once and reused everywhere it is needed downstream.

When the graph has multiple paths from a parameter to the loss, the gradients along each path **add**. This is why a shared embedding table — used a hundred times in a transformer — ends up with the sum of contributions from all hundred uses. You get this for free; the recursion described below visits every path naturally.

Consider this computation:  $\text{loss} = (x * W + b)^2$

Forward:  $x \rightarrow [\text{Mul}(W)] \rightarrow z_1 \rightarrow [\text{Add}(b)] \rightarrow z_2 \rightarrow [\text{Square}] \rightarrow \text{loss}$

Backward chain rule:

|                                                                                 |                           |
|---------------------------------------------------------------------------------|---------------------------|
| $\partial \text{loss} / \partial z_2 = 2 * z_2$                                 | (square backward)         |
| $\partial \text{loss} / \partial z_1 = \partial \text{loss} / \partial z_2 * 1$ | (addition backward)       |
| $\partial \text{loss} / \partial x = \partial \text{loss} / \partial z_1 * W$   | (multiplication backward) |

Each backward function does exactly one thing: multiply the incoming gradient by its own local derivative. It knows nothing about the rest of the graph — and it doesn't need to. Here's how `MulBackward` implements this:

The code in `?lst-06-autograd-mul-backward` makes this concrete.

```
class MulBackward(Function):
    """Gradient computation for element-wise multiplication."""

    def apply(self, grad_output):
        """
        For  $z = a * b$ :
         $\partial z / \partial a = b \rightarrow \text{grad\_a} = \text{grad\_output} * b$ 
         $\partial z / \partial b = a \rightarrow \text{grad\_b} = \text{grad\_output} * a$ 

        Uses vectorized element-wise multiplication (NumPy broadcasting).
        """
        a, b = self.saved_tensors
        grad_a = grad_b = None

        if a.requires_grad:
            grad_a = grad_output * b.data # Vectorized element-wise multiplication

        if b.requires_grad:
            grad_b = grad_output * a.data # NumPy handles broadcasting automatically

        return grad_a, grad_b
```

: Listing 6.2 — `MulBackward` applies the product rule with NumPy broadcasting instead of explicit loops. {`lst-06-autograd-mul-backward`}

Each operation knows only its own derivative; the chain rule does the connecting. NumPy handles the element-wise math in optimized C, so no explicit loops are needed.

### 10.5.3 Backward Pass Implementation

The backward pass walks the computation graph in reverse, computing gradients for every tensor it visits. Your `backward()` method does this as a recursive tree walk — short enough to read in one sitting, but enough to support arbitrarily deep networks:

The code in `?lst-06-autograd-tensor-backward` makes this concrete.

```
def backward(self, gradient=None):
    """Compute gradients via backpropagation."""
    if not self.requires_grad:
        return

    # Initialize gradient for scalar outputs
    if gradient is None:
        if self.data.size == 1:
            gradient = np.ones_like(self.data)
        else:
```

```

        raise ValueError("backward() requires gradient for non-scalar tensors")

    # Accumulate gradient (vectorized NumPy operation)
    if self.grad is None:
        self.grad = np.zeros_like(self.data)
    self.grad += gradient

    # Propagate to parent tensors
    if hasattr(self, '_grad_fn') and self._grad_fn is not None:
        grads = self._grad_fn.apply(gradient) # Compute input gradients using
        vectorized ops

    for tensor, grad in zip(self._grad_fn.saved_tensors, grads):
        if isinstance(tensor, Tensor) and tensor.requires_grad and grad is not
        None:
            tensor.backward(grad) # Recursive call

```

: Listing 6.3 — `Tensor.backward()` seeds the output gradient, accumulates into `.grad`, and recurses through the `_grad_fn` chain. {#lst-06-autograd-tensor-backward}

For a 100-layer network, `loss.backward()` triggers 100 recursive calls — one per layer — flowing gradients from output to input. The traversal is recursive Python; the math inside each `apply()` is vectorized NumPy. That split is why the system stays both readable and fast.

The gradient argument deserves a closer look. For scalar losses (the typical case) you call `loss.backward()` with no arguments and the method seeds the gradient to 1.0 — because  $\partial_{\text{loss}}/\partial_{\text{loss}} = 1$ . For non-scalar outputs you must pass the upstream gradient explicitly; there is no canonical scalar to seed from, and silently picking one would hide bugs.

### 10.5.4 Gradient Accumulation

Gradient accumulation is the same feature seen from two sides. Call `backward()` twice on the same tensor and the gradients add — by design. That’s what lets you split a batch that doesn’t fit in memory into chunks, run them sequentially, and end up with the same gradient as if you’d processed them all at once:

```

# Large batch (doesn't fit in memory)
for mini_batch in split_batch(large_batch, chunks=4):
    loss = model(mini_batch)
    loss.backward() # Gradients accumulate in model parameters

# Now gradients equal the sum over the entire large batch
optimizer.step()
model.zero_grad() # Reset for next iteration

```

Without this behavior you’d have to store every mini-batch gradient and sum them yourself. With it, the autograd system does the bookkeeping.

The flip side: accumulation becomes a silent bug the moment you forget to call `zero_grad()` between iterations.

```

# WRONG: Gradients accumulate across iterations
for batch in dataloader:

```

```

    loss = model(batch)
    loss.backward() # Gradients keep adding!
    optimizer.step() # Updates use accumulated gradients from all previous batches

# CORRECT: Zero gradients after each update
for batch in dataloader:
    model.zero_grad() # Reset gradients
    loss = model(batch)
    loss.backward()
    optimizer.step()

```

Your `zero_grad()` implementation is simple but crucial:

```

def zero_grad(self):
    """Reset gradients to None."""
    self.grad = None

```

Setting to `None` instead of zeros saves memory: NumPy doesn't allocate arrays until you accumulate the first gradient.

### 10.5.5 Memory Management in Autograd

Autograd's memory footprint comes from two sources: stored intermediate tensors and gradient storage. For a forward pass through an  $N$ -layer network, you store roughly  $N$  intermediate activations. During backward pass, you store gradients for every parameter.

Consider a simple linear layer:  $y = x @ W + b$

**Forward pass stores:** -  $x$  (needed for computing  $\text{grad}_W = x.T @ \text{grad}_y$ ) -  $W$  (needed for computing  $\text{grad}_x = \text{grad}_y @ W.T$ )

**Backward pass allocates:** -  $\text{grad}_x$  (same shape as  $x$ ) -  $\text{grad}_W$  (same shape as  $W$ ) -  $\text{grad}_b$  (same shape as  $b$ )

For a batch of 32 samples through a (512, 768) linear layer, the memory breakdown is:

```

Forward storage:
x:      32 × 512 × 4 bytes =      64 KB
W:      512 × 768 × 4 bytes =  1,536 KB

Backward storage:
grad_x:  32 × 512 × 4 bytes =      64 KB
grad_W: 512 × 768 × 4 bytes =  1,536 KB
grad_b:   768 × 4 bytes =         3 KB

Total: ~3.1 MB for one layer (2× parameter size + activation size)

```

Multiply by network depth and you see why memory limits batch size. A 100-layer transformer stores  $100\times$  the activations, which can easily exceed GPU memory.

Production frameworks mitigate this with gradient checkpointing: they discard intermediate activations during forward pass and recompute them during backward pass. This trades compute (recomputing activations) for memory (not storing them). Your implementation doesn't do this—it's an advanced optimization—but understanding the trade-off is essential.

The implementation shows this memory overhead clearly in the `MatmulBackward` class:

The code in `lst-06-autograd-matmul-backward` makes this concrete.

```
class MatmulBackward(Function):
    """
    Gradient computation for matrix multiplication.

    For  $Z = A @ B$ :
    - Must store A and B during forward pass
    - Backward computes:  $\text{grad}_A = \text{grad}_Z @ B.T$  and  $\text{grad}_B = A.T @ \text{grad}_Z$ 
    - Uses vectorized NumPy operations (np.matmul, np.swapaxes)
    """

    def apply(self, grad_output):
        a, b = self.saved_tensors  # Retrieved from memory
        grad_a = grad_b = None

        if a.requires_grad:
            # Vectorized transpose and matmul (no explicit loops)
            b_T = np.swapaxes(b.data, -2, -1)
            grad_a = np.matmul(grad_output, b_T)

        if b.requires_grad:
            # Vectorized operations for efficiency
            a_T = np.swapaxes(a.data, -2, -1)
            grad_b = np.matmul(a_T, grad_output)

        return grad_a, grad_b
```

: Listing 6.4 — `MatmulBackward` turns one forward `matmul` into two backward `matmuls`, the source of the 2x backward-to-forward FLOP ratio. `{#lst-06-autograd-matmul-backward}`

Notice that both `a` and `b` must be saved during forward pass. For large matrices, this storage cost dominates memory usage. All gradient computations use vectorized NumPy operations, which are implemented in optimized C/Fortran code under the hood—no explicit Python loops are needed.

## 10.6 Production Context

### 10.6.1 Your Implementation vs. PyTorch

Your autograd system and PyTorch’s share the same design: computation graphs built during forward pass, reverse-mode differentiation during backward pass, and gradient accumulation in parameter tensors. The differences are in scale and optimization.

Table 10.5 places your implementation side by side with the production reference for direct comparison.

Table 10.5: Feature comparison between TinyTorch autograd and PyTorch autograd.

| Feature               | Your Implementation                             | PyTorch                     |
|-----------------------|-------------------------------------------------|-----------------------------|
| <b>Graph Building</b> | Python objects, <code>_grad_fn</code> attribute | C++ objects, compiled graph |

| Feature           | Your Implementation        | PyTorch                                                    |
|-------------------|----------------------------|------------------------------------------------------------|
| <b>Memory</b>     | Stores all intermediates   | Gradient checkpointing, memory pools                       |
| <b>Speed</b>      | Pure Python, NumPy backend | C++/CUDA, fused kernels                                    |
| <b>Operations</b> | 10 backward functions      | 2000+ optimized backward functions                         |
| <b>Debugging</b>  | Direct Python inspection   | <code>torch.autograd.profiler</code> , graph visualization |

### 10.6.2 Code Comparison

The following comparison shows identical conceptual patterns in TinyTorch and PyTorch. The APIs mirror each other because both implement the same autograd algorithm.

## 10.7 Your TinyTorch

```
from tinytorch import Tensor

# Create tensors with gradient tracking
x = Tensor([[1.0, 2.0]], requires_grad=True)
W = Tensor([[3.0], [4.0]], requires_grad=True)

# Forward pass builds computation graph
y = x.matmul(W) # y = x @ W
loss = (y * y).sum() # loss = sum(y^2)

# Backward pass computes gradients
loss.backward()

# Access gradients
print(f"x.grad: {x.grad}") # ∂loss/∂x
print(f"W.grad: {W.grad}") # ∂loss/∂W
```

## 10.8 PyTorch

```
import torch

# Create tensors with gradient tracking
x = torch.tensor([[1.0, 2.0]], requires_grad=True)
W = torch.tensor([[3.0], [4.0]], requires_grad=True)

# Forward pass builds computation graph
y = x @ W # PyTorch uses @ operator
```



```

loss = (y * y).sum()

# Backward pass computes gradients
loss.backward()

# Access gradients
print(f"x.grad: {x.grad}")
print(f"W.grad: {W.grad}")

```

Let's walk through the comparison line by line:

- **Line 3-4 (Tensor creation):** Both frameworks use `requires_grad=True` to enable gradient tracking. This is an opt-in design: most tensors (data, labels) don't need gradients, only parameters do.
- **Line 7-8 (Forward pass):** Operations automatically build computation graphs. TinyTorch uses `.matmul()` method; PyTorch supports both `.matmul()` and the `@` operator.
- **Line 11 (Backward pass):** Single method call triggers reverse-mode differentiation through the entire graph.
- **Line 14-15 (Gradient access):** Both store gradients in the `.grad` attribute. Gradients have the same shape as the original tensor.

#### 💡 What's Identical

Computation graph construction, chain rule implementation, and gradient accumulation semantics. When you debug PyTorch autograd issues, you're debugging the same algorithm you implemented here.

### 10.8.1 Why Autograd Matters at Scale

The case for automating differentiation is overwhelming the moment you look at the numbers:

- **GPT-3:** 175 billion parameters — that's 175,000,000,000 gradients per training step.
- **Training cost:** each backward pass takes roughly  $2\times$  the forward pass time (two matmuls instead of one per linear layer).
- **Memory:** storing the computation graph for a transformer can require  $\sim 10\times$  the model's parameter footprint.

Hand-deriving gradients does not scale to any of this. Even a 3-layer MLP with a million parameters would take weeks to differentiate manually and would still contain bugs at the end. Autograd makes training tractable by automating the most error-prone part of deep learning — and it's the single biggest reason the field moves as fast as it does.

## 10.9 Check Your Understanding

### 💡 Check Your Understanding — Autograd

Before moving on, verify you can articulate each of the following:

- Why activations must be pinned in memory during the forward pass and can only be freed after the matching backward step completes.
- Why the backward pass costs roughly  $2\times$  the forward pass FLOPs (one extra matmul per linear layer for `grad_x` and `grad_w`).

- How forgetting `zero_grad()` silently accumulates gradients across iterations and what that does to update magnitude and direction.
- Why training a 1B-parameter model with Adam needs ~16 GB of framework overhead before storing a single activation — and what that implies for batch size.

If any of these feels fuzzy, revisit the Computation Graphs and Memory Management sections before moving on.

Test yourself with these systems thinking questions. They're designed to build intuition for autograd's performance characteristics and design decisions.

### Q1: Computation Graph Memory

A 5-layer MLP processes a batch of 64 samples. Each layer stores its input activation for backward pass. Layer dimensions are:  $784 \rightarrow 512 \rightarrow 256 \rightarrow 128 \rightarrow 10$ . How much memory (in MB) is used to store activations for one batch?

#### 💡 Answer

```
Layer 1 input: 64 × 784 × 4 bytes = 196 KB
Layer 2 input: 64 × 512 × 4 bytes = 128 KB
Layer 3 input: 64 × 256 × 4 bytes = 64 KB
Layer 4 input: 64 × 128 × 4 bytes = 32 KB
Layer 5 input: 64 × 10 × 4 bytes = 2.5 KB
```

**Total: ~422 KB  $\approx$  0.41 MB**

That is per forward pass through a tiny MLP. A 100-layer transformer stores roughly  $100\times$  this — and with much wider layers — which is why gradient checkpointing trades compute for memory by recomputing activations during backward pass.

### Q2: Backward Pass Complexity

A forward pass through a linear layer  $y = x @ W$  (where  $x$  is  $32 \times 512$  and  $W$  is  $512 \times 256$ ) takes 8ms. How long will the backward pass take?

#### 💡 Answer

Forward: 1 matmul ( $x @ W$ )

Backward: 2 matmuls -  $\text{grad}_x = \text{grad}_y @ W.T$  ( $32 \times 256 @ 256 \times 512$ ) -  $\text{grad}_W = x.T @ \text{grad}_y$  ( $512 \times 32 @ 32 \times 256$ )

**Backward takes  $\sim 2\times$  forward time  $\approx$  16ms**

This is why training (forward + backward) takes roughly  $3\times$  inference time. GPU parallelism and kernel fusion can reduce this, but the fundamental 2:1 ratio remains.

### Q3: Gradient Accumulation Memory

You have 16GB GPU memory and a model with 1B parameters (float32). How much memory is available for activations and gradients during training?

 Answer

```

Model parameters:      1B × 4 bytes = 4 GB
Gradients:              1B × 4 bytes = 4 GB
Optimizer state (Adam): 1B × 8 bytes = 8 GB    (momentum + variance)

```

**Total framework overhead: 16 GB**

**Available for activations: 0 GB** — you’ve already exceeded memory before storing a single activation. This is why large models reach for gradient accumulation across multiple forward passes before updating parameters, or gradient checkpointing to shrink activation memory. The “2× parameter size” rule (params + grads) is a floor, not a ceiling — optimizers like Adam add more on top.

**Q4: requires\_grad Performance**

A typical training batch has: 32 images (input), 10M parameter tensors (weights), 50 intermediate activation tensors. If `requires_grad` defaults to `True` for all tensors, how many tensors unnecessarily track gradients?

 Answer

Tensors that **need** gradients:

- Parameters: 10M tensors

Tensors that **don’t** need gradients:

- Input images: 32 tensors (data is not learned)
- Intermediate activations: 50 tensors (needed for backward, but not updated)

**32 input tensors would unnecessarily track gradients** if `requires_grad` defaulted to `True`.

This is why PyTorch defaults `requires_grad=False` for new tensors and forces an explicit opt-in for parameters. For a batch of 32 images of shape  $3 \times 224 \times 224$ , accidentally tracking gradients on the inputs wastes  $4.8\text{M float32 values} \times 4 \text{ bytes} \approx 18.4 \text{ MB}$  per batch — for nothing.

**Q5: Graph Retention**

You forget to call `zero_grad()` before each training iteration. After 10 iterations, how do the gradients compare to correct training?

 Answer

**Gradients accumulate across all 10 iterations.**

If correct gradient for iteration  $i$  is  $g_i$ , your accumulated gradient is:  $\text{grad} = g_1 + g_2 + g_3 + \dots + g_{10}$

**Effects:** 1. **Magnitude:** Gradients are  $\sim 10\times$  larger than they should be 2. **Direction:** The sum of 10 different gradients, which may not point toward the loss minimum 3. **Learning:** Parameter updates use the wrong direction and wrong magnitude 4. **Result:** Training diverges or oscillates instead of converging

**Bottom line:** Always call `zero_grad()` at the start of each iteration (or after `optimizer.step()`).

## 10.10 Key Takeaways

- **Computation graphs are dynamic and implicit:** Every tensor produced by an operation stores a `_grad_fn` pointer; that pointer chain *is* the graph autograd walks in reverse.

- **Activation pinning is autograd’s memory tax:** Saved tensors survive until `backward()` runs, which is why forward-pass VRAM usage scales linearly with depth and caps batch size long before parameter count does.
- **Backward costs  $2\times$  forward:** Each linear layer turns one matmul into two (`grad_x`, `grad_w`), making training roughly  $3\times$  the cost of inference — a constant every capacity plan assumes.
- **`zero_grad()` is not optional:** Gradient accumulation is the feature *and* the silent bug; skipping the reset compounds 10 gradients into one bad step.

**Coming next:** You can compute a gradient for every parameter — but a gradient alone is just a direction. Module 07 turns directions into updates by building SGD, Adam, and AdamW, the rules that decide how far to step.

## 10.11 Further Reading

For students who want to understand the academic foundations and mathematical underpinnings of automatic differentiation:

### 10.11.1 Seminal Papers

- **Automatic Differentiation in Machine Learning: a Survey** - Baydin et al. (2018). Comprehensive survey of AD techniques, covering forward-mode, reverse-mode, and mixed-mode differentiation. Essential reading for understanding autograd theory.
  - **Systems Implication:** This work formalized the foundational memory-compute trade-off: reverse-mode AD minimizes compute at the steep cost of  $O(N)$  memory to store the forward activation graph, permanently dictating modern GPU VRAM requirements. [arXiv:1502.05767](#)
- **Automatic Differentiation of Algorithms** - Griewank (1989). The foundational work on reverse-mode AD that underlies all modern deep learning frameworks. Introduces the mathematical formalism for gradient computation via the chain rule.
  - **Systems Implication:** Exposed the inherent sequential data dependencies during the backward pass. Because gradients must be propagated backwards sequentially from output to input, this creates a structural bottleneck that prevents perfect parallelization across network layers. [Computational Optimization and Applications](#)
- **PyTorch: An Imperative Style, High-Performance Deep Learning Library** - Paszke et al. (2019). Describes PyTorch’s autograd implementation and design philosophy. Shows how imperative programming (define-by-run) enables dynamic computation graphs.
  - **Systems Implication:** By introducing dynamic computation graphs, PyTorch mandated that memory allocation and operator dispatch overhead occur strictly at runtime (eager execution). This paradigm shift required highly optimized C++ backends to aggressively hide Python’s interpretive latency. [NeurIPS 2019](#)

### 10.11.2 Additional Resources

- **Textbook:** [“Deep Learning”](#) by Goodfellow, Bengio, and Courville - Chapter 6 covers backpropagation and computational graphs with excellent visualizations
- **Tutorial:** [CS231n: Backpropagation, Intuitions](#) - Stanford’s visual explanation of gradient flow through computation graphs
- **Documentation:** [PyTorch Autograd Mechanics](#) - Official guide to PyTorch’s autograd implementation details

## 10.12 What's Next

You can now compute a gradient for every parameter in any network you build. That gradient tells you which way is downhill — but it does not tell you how big a step to take, or how to dampen oscillations, or how to adapt the step size per-parameter. That is the optimizer's job.

### Coming Up: Module 07 — Optimizers

You'll implement SGD, momentum, and Adam: the rules that turn the `param.grad` tensors produced by `backward()` into actual parameter updates. With autograd plus an optimizer, you have the entire machinery a training loop needs.

**Preview — how your autograd gets used in the modules that follow:**

Table 10.6 traces how this module is reused by later parts of the curriculum.

Table 10.6: How autograd feeds into subsequent optimizer and training modules.

| Module                | What It Does                      | Your Autograd In Action                                                                        |
|-----------------------|-----------------------------------|------------------------------------------------------------------------------------------------|
| <b>07: Optimizers</b> | Update parameters using gradients | <code>optimizer.step()</code> uses <code>param.grad</code> computed by <code>backward()</code> |
| <b>08: Training</b>   | Complete training loops           | <code>loss.backward()</code> → <code>optimizer.step()</code> → repeat                          |
| <b>12: Attention</b>  | Multi-head self-attention         | Gradients flow through Q, K, V projections automatically                                       |

## 10.13 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.

ne, or clone the repository for persistent local work. ::: ocal work. :::



## 🔥 Chapter 11

# Module 07: Optimizers

A gradient is a direction, not a step. The optimizer picks a step size, a history window, and a per-parameter scaling rule. Its state, momentum buffers, running second moments, copies of every weight, often costs more memory than the model itself. Choose well and a model converges in an afternoon. Choose poorly and VRAM runs out on epoch one.

### Module Info

**FOUNDATION TIER** | Difficulty: ●●○○ | Time: 3-5 hours | Prerequisites: 01-06

**Prerequisites:** Modules 01-06 means you need:

- Tensor operations and parameter storage
- DataLoader for efficient batch processing
- Understanding of forward/backward passes (autograd)
- Why gradients point toward higher loss

If you understand how `loss.backward()` computes gradients and why we need to update parameters to minimize loss, you're ready.

## 11.1 Overview

You have gradients. Now what? An optimizer is the rule that turns a gradient into a parameter update — the difference between a model that converges in an afternoon and one that diverges on the first batch. Picture optimization as hiking in fog: you can feel the slope under your feet but cannot see the valley. Each optimizer is a different strategy for choosing your next step.

You'll build three: SGD with momentum (the foundation), Adam with adaptive per-parameter learning rates (the modern workhorse), and AdamW with decoupled weight decay (the default for transformers). They differ in memory cost, convergence speed, and how forgiving they are when you guess the learning rate wrong — and that last point matters more than most practitioners admit.

## 11.2 Learning Objectives

### By completing this module, you will:

- **Implement** SGD with momentum to reduce oscillations and accelerate convergence in narrow valleys
- **Master** Adam's adaptive learning rate mechanism with first and second moment estimation
- **Understand** memory trade-offs (SGD: 2x memory vs Adam: 3x memory) and computational complexity per step
- **Connect** optimizer state management to checkpointing and distributed training considerations

## 11.3 What You'll Build




Figure 11.1: **TinyTorch Optimizer Hierarchy:** From basic SGD to advanced adaptive algorithms.

### Implementation roadmap:

Table 11.1 lays out the implementation in order, one part at a time.

Table 11.1: **Implementation roadmap for the SGD and Adam optimizers.**

| Step | What You'll Implement | Key Concept                                                         |
|------|-----------------------|---------------------------------------------------------------------|
| 1    | Optimizer base class  | Common interface:<br><code>zero_grad()</code> , <code>step()</code> |
| 2    | SGD with momentum     | Velocity buffers to reduce oscillations                             |
| 3    | Adam optimizer        | First and second moment estimation with bias correction             |
| 4    | AdamW optimizer       | Decoupled weight decay for proper regularization                    |

### The pattern you'll enable:

```

# Training loop with optimizer
optimizer = Adam(model.parameters(), lr=0.001)
loss.backward() # Compute gradients (Module 06)
optimizer.step() # Update parameters using gradients
optimizer.zero_grad() # Clear gradients for next iteration
  
```

### 11.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Learning rate schedules (that's Module 08: Training)
- Gradient clipping (PyTorch provides this via `torch.nn.utils.clip_grad_norm_`)
- Second-order optimizers like L-BFGS (rarely used in deep learning due to memory cost)
- Distributed optimizer sharding (production frameworks use techniques like ZeRO)



**You are building the core optimization algorithms.** Advanced training techniques come in Module 08.

## 11.4 API Reference

The signatures your implementations must satisfy. Keep this open while you build.

### 11.4.1 Optimizer Base Class

```
Optimizer(params: List[Tensor])
```

Base class defining the optimizer interface. All optimizers inherit from this.  
Table 11.2 lists the base class surface.

Table 11.2: **Methods defined by the Optimizer base class.**

| Method    | Signature           | Description                                   |
|-----------|---------------------|-----------------------------------------------|
| zero_grad | zero_grad() -> None | Clear gradients from all parameters           |
| step      | step() -> None      | Update parameters (implemented by subclasses) |

### 11.4.2 SGD Optimizer

```
SGD(params, lr=0.01, momentum=0.0, weight_decay=0.0)
```

Stochastic Gradient Descent with optional momentum and weight decay.

**Parameters:** - `params`: List of Tensor parameters to optimize - `lr`: Learning rate (step size, default: 0.01)  
- `momentum`: Momentum factor (0.0-1.0, typically 0.9, default: 0.0) - `weight_decay`: L2 penalty coefficient (default: 0.0)

**Update rule:** - Without momentum:  $\text{param} = \text{param} - \text{lr} * \text{grad}$  - With momentum:  $v = \text{momentum} * v + \text{grad}$ ;  $\text{param} = \text{param} - \text{lr} * v$

**State management methods:**

Table 11.3 lists the state-management methods.

Table 11.3: **State-management methods on the SGD optimizer.**

| Method             | Signature                                         | Description                                     |
|--------------------|---------------------------------------------------|-------------------------------------------------|
| has_momentum       | has_momentum() -> bool                            | Check if optimizer uses momentum (momentum > 0) |
| get_momentum_state | get_momentum_state() -> Optional[List]            | Get momentum buffers for checkpointing          |
| set_momentum_state | set_momentum_state(state: Optional[List]) -> None | Restore momentum buffers from checkpoint        |

### 11.4.3 Adam Optimizer

```
Adam(params, lr=0.001, betas=(0.9, 0.999), eps=1e-8, weight_decay=0.0)
```

Adaptive Moment Estimation with per-parameter learning rates.

**Parameters:** - `params`: List of Tensor parameters to optimize - `lr`: Learning rate (default: 0.001) - `betas`: Tuple of coefficients ( $\beta_1, \beta_2$ ) for computing running averages (default: (0.9, 0.999)) - `eps`: Small constant for numerical stability (default: 1e-8) - `weight_decay`: L2 penalty coefficient (default: 0.0)

**State:** - `m_buffers`: First moment estimates (momentum of gradients) - `v_buffers`: Second moment estimates (momentum of squared gradients)

### 11.4.4 AdamW Optimizer

```
AdamW(params, lr=0.001, betas=(0.9, 0.999), eps=1e-8, weight_decay=0.01)
```

Adam with decoupled weight decay regularization.

**Parameters:** - `params`: List of Tensor parameters to optimize - `lr`: Learning rate (default: 0.001) - `betas`: Tuple of coefficients ( $\beta_1, \beta_2$ ) for computing running averages (default: (0.9, 0.999)) - `eps`: Small constant for numerical stability (default: 1e-8) - `weight_decay`: L2 penalty coefficient (default: 0.01, higher than Adam)

**Key difference from Adam:** Weight decay is applied directly to parameters after gradient update, not mixed into the gradient.

## 11.5 Core Concepts

Four ideas carry the rest of this chapter: the descent rule itself, why momentum exists, why Adam adapts per-parameter, and why the learning rate dwarfs every other knob.

### 11.5.1 Gradient Descent Fundamentals

Gradient descent is conceptually simple: gradients point uphill toward higher loss, so we step downhill by moving in the opposite direction. The gradient  $\nabla L$  tells us the direction of steepest ascent, so  $-\nabla L$  points toward steepest descent.

The basic update rule is:  $\theta_{\text{new}} = \theta_{\text{old}} - \alpha * \nabla L$ , where  $\theta$  represents parameters and  $\alpha$  is the learning rate (step size). This simple formula hides important challenges. How large should steps be? What if different parameters need different step sizes? What about noisy gradients or narrow valleys that cause oscillation?

Here's how your SGD implementation handles the basic case without momentum:

The code in ?@lst-07-optimizers-sgd-step makes this concrete.

```
def step(self):
    """Perform SGD update step with momentum."""
    for i, param in enumerate(self.params):
        if param.grad is None:
            continue

        # Get gradient data
        grad = param.grad
        if isinstance(grad, Tensor):
            grad_data = grad.data
        else:
```

```

        grad_data = grad

        # Apply weight decay if specified
        if self.weight_decay != 0:
            grad_data = grad_data + self.weight_decay * param.data

        # Update parameter: param = param - lr * grad
        param.data = param.data - self.lr * grad_data

    self.step_count += 1

```

: Listing 7.1 — Plain SGD `step()`: the one-line update `param = param - lr * grad` wrapped in parameter iteration and optional weight decay. {#lst-07-optimizers-sgd-step}

The code reveals the simplicity of basic SGD: subtract learning rate times gradient from each parameter. But this simplicity comes with a cost: plain SGD can oscillate wildly in narrow valleys of the loss landscape.

**Complexity:** SGD’s update is  $O(P)$  compute and  $O(P)$  memory per step, where  $P$  is the parameter count; adding momentum costs an additional  $O(P)$  memory for the velocity buffer. Adam is also  $O(P)$  compute per step but pays  $O(2P)$  extra memory for the first and second moment buffers ( $m$  and  $v$ ). The asymptotic cost of every optimizer in this chapter scales *linearly* with parameter count — what differs is the constant factor in memory.

### 11.5.2 Momentum and Acceleration

Momentum solves the oscillation problem by remembering previous update directions. Think of a ball rolling down a hill: it doesn’t immediately change direction when it hits a small bump because it has momentum carrying it forward. In optimization, momentum accumulates velocity in directions that gradients consistently agree on, while oscillations in perpendicular directions cancel out.

The momentum update maintains a velocity buffer  $v$  for each parameter:  $v = \beta * v_{\text{prev}} + \text{grad}$  and then `param = param - lr * v`. The momentum coefficient  $\beta$  (typically 0.9) controls how much previous direction we remember. With  $\beta=0.9$ , we keep 90% of the old velocity and add 10% of the current gradient.

Here’s how your SGD implementation adds momentum:

The code in ?@lst-07-optimizers-sgd-momentum makes this concrete.

```

# Update momentum buffer
if self.momentum != 0:
    if self.momentum_buffers[i] is None:
        # Initialize momentum buffer on first use
        self.momentum_buffers[i] = np.zeros_like(param.data)

    # Update momentum: v = momentum * v_prev + grad
    self.momentum_buffers[i] = self.momentum * self.momentum_buffers[i] + grad_data
    grad_data = self.momentum_buffers[i]

# Update parameter: param = param - lr * grad
param.data = param.data - self.lr * grad_data

```

: Listing 7.2 — SGD momentum block: lazy buffer initialisation plus the exponential-moving-average update `v = beta * v + grad`. {#lst-07-optimizers-sgd-momentum}

The momentum buffer is initialized lazily (only when first needed) to save memory for optimizers without momentum. Once initialized, each step accumulates 90% of the previous velocity plus the current gradient, creating a smoothed update direction that's less susceptible to noise and oscillation.

### 11.5.3 Adam and Adaptive Learning Rates

Different parameters want different learning rates. An embedding weight that lives in  $[-0.01, 0.01]$  and an output weight that lives in  $[-10, 10]$  cannot share a step size — one rate is too small for the embedding, too large for the output. SGD ignores this. Adam fixes it.

Adam keeps two running averages per parameter: a first moment  $m$  (mean of gradients) and a second moment  $v$  (mean of squared gradients). Updating with  $m / \sqrt{v}$  gives an automatic per-parameter step size — parameters with consistently large gradients take smaller steps; parameters with small or noisy gradients take larger ones.

The algorithm tracks:  $m = \beta_1 * m_{\text{prev}} + (1 - \beta_1) * \text{grad}$  and  $v = \beta_2 * v_{\text{prev}} + (1 - \beta_2) * \text{grad}^2$ . Then it corrects for initialization bias ( $m$  and  $v$  start at zero) and updates:  $\text{param} = \text{param} - \text{lr} * \hat{m} / (\sqrt{\hat{v}} + \epsilon)$ , where  $\hat{m}$  and  $\hat{v}$  are bias-corrected moments.

Here's the complete Adam update from your implementation:

The code in `?@lst-07-optimizers-adam-step` makes this concrete.

```
def step(self):
    """Perform Adam update step."""
    self.step_count += 1

    for i, param in enumerate(self.params):
        if param.grad is None:
            continue

        grad = param.grad
        if isinstance(grad, Tensor):
            grad_data = grad.data
        else:
            grad_data = grad

        # Initialize buffers if needed
        if self.m_buffers[i] is None:
            self.m_buffers[i] = np.zeros_like(param.data)
            self.v_buffers[i] = np.zeros_like(param.data)

        # Update biased first moment estimate
        self.m_buffers[i] = self.beta1 * self.m_buffers[i] + (1 - self.beta1) * grad_data

        # Update biased second moment estimate
        self.v_buffers[i] = self.beta2 * self.v_buffers[i] + (1 - self.beta2) * (grad_data ** 2)

        # Compute bias correction
        bias_correction1 = 1 - self.beta1 ** self.step_count
        bias_correction2 = 1 - self.beta2 ** self.step_count
```

```

# Compute bias-corrected moments
m_hat = self.m_buffers[i] / bias_correction1
v_hat = self.v_buffers[i] / bias_correction2

# Update parameter
param.data = param.data - self.lr * m_hat / (np.sqrt(v_hat) + self.eps)

```

: Listing 7.3 — Full Adam `step()`: first and second moment updates, bias correction, and the per-parameter adaptive step  $\text{lr} * \text{m\_hat} / (\text{sqrt}(\text{v\_hat}) + \text{eps})$ . {#lst-07-optimizers-adam-step}

The bias correction terms  $(1 - \beta^t)$  are crucial in the first few steps. Without correction,  $m$  and  $v$  start at zero and take many steps to reach reasonable values, causing the optimizer to take tiny steps initially. The correction divides by increasingly large values: at step 1, divide by 0.1; at step 2, divide by 0.19; eventually the correction approaches 1.0 and has no effect.

While Adam’s adaptive step sizes dramatically accelerate convergence, this mathematical elegance imposes a severe, often debilitating, hardware penalty on large-scale training systems.

#### **i** Systems Implication: Optimizer State Memory

Because Adam maintains both a first moment (`m_buffers`) and a second moment (`v_buffers`) for every individual parameter, its memory footprint is enormous. Training a 1GB model requires 1GB for parameters, 1GB for gradients, and an additional **2GB solely for Adam’s state**—meaning the optimizer demands  $3\times$  the model’s footprint in VRAM. Furthermore, continuously fetching these massive buffers from GPU memory to compute each update makes the optimizer step heavily **memory-bound**, bottlenecked fundamentally by VRAM memory bandwidth rather than streaming multiprocessor compute cores.

### 11.5.4 AdamW and Decoupled Weight Decay

AdamW fixes a real bug in Adam. Standard Adam folds weight decay into the gradient —  $\text{grad} = \text{grad} + \lambda * \text{param}$  — and then runs the adaptive update. The decay term then gets divided by  $\sqrt{v}$  along with everything else, so parameters with large gradients receive *less* regularization and parameters with small gradients receive *more*. That is the opposite of what you want.

AdamW decouples the two: take the normal Adam step using the raw gradient, then separately shrink each parameter by a fixed fraction  $\text{lr} * \text{weight\_decay}$ . Regularization strength is now constant across the network, independent of gradient magnitude. This is why every modern transformer trains with AdamW, not Adam.

Here’s how your AdamW implementation achieves decoupling:

The code in {#lst-07-optimizers-adamw-step} makes this concrete.

```

# Update moments using pure gradients (NO weight decay mixed in)
self.m_buffers[i] = self.beta1 * self.m_buffers[i] + (1 - self.beta1) * grad_data
self.v_buffers[i] = self.beta2 * self.v_buffers[i] + (1 - self.beta2) * (grad_data **
    2)

# Compute bias correction and bias-corrected moments
bias_correction1 = 1 - self.beta1 ** self.step_count
bias_correction2 = 1 - self.beta2 ** self.step_count

```

```

m_hat = self.m_buffers[i] / bias_correction1
v_hat = self.v_buffers[i] / bias_correction2

# Apply gradient-based update
param.data = param.data - self.lr * m_hat / (np.sqrt(v_hat) + self.eps)

# Apply decoupled weight decay (separate from gradient update)
if self.weight_decay != 0:
    param.data = param.data * (1 - self.lr * self.weight_decay)

```

: Listing 7.4 — AdamW update: the adaptive Adam step runs on pure gradients, then weight decay shrinks parameters separately by  $(1 - \text{lr} * \text{weight\_decay})$ . {#lst-07-optimizers-adamw-step}

Notice that weight decay appears only at the end, multiplying parameters by  $(1 - \text{lr} * \text{weight\_decay})$  to shrink them slightly. This shrinkage happens after the gradient update and is completely independent of gradient magnitudes or adaptive scaling.

### 11.5.5 Learning Rate Selection

The learning rate is the most important hyperparameter you will ever set. Too large and parameters oscillate or diverge; too small and training stalls or settles into a bad minimum. No optimizer choice — not Adam, not AdamW — rescues a badly chosen learning rate.

SGD typically lives in  $0.001$  to  $0.1$  and is unforgiving: get it wrong by 10x and training breaks. Momentum smooths the ride but does not change the sensitivity. Adam and AdamW default to  $0.001$  and tolerate roughly an order of magnitude in either direction, which is a large part of why they are popular. Transformers are the exception — they want  $1e-4$  to  $3e-4$  with a warmup that ramps the rate from zero over the first few thousand steps.

Batch size couples to learning rate. Larger batches produce less noisy gradients and can absorb larger steps. The standard heuristic — scale the learning rate linearly with batch size — works up to a few thousand samples per batch and breaks beyond that, which is why large-scale training has its own literature.

## 11.6 Production Context

### 11.6.1 Your Implementation vs. PyTorch

Your TinyTorch optimizers and PyTorch’s `torch.optim` share the same algorithmic foundations and API patterns. The differences lie in implementation details: PyTorch uses optimized C++/CUDA kernels, supports mixed precision training, and includes specialized optimizers for specific domains.

Table 11.4 places your implementation side by side with the production reference for direct comparison.

Table 11.4: Feature comparison between TinyTorch optimizers and `torch.optim`.

| Feature                 | Your Implementation  | PyTorch                             |
|-------------------------|----------------------|-------------------------------------|
| <b>Backend</b>          | NumPy (Python)       | C++/CUDA kernels                    |
| <b>Speed</b>            | 1x (baseline)        | 10-50x faster                       |
| <b>Memory</b>           | Same asymptotic cost | Same (3x for Adam)                  |
| <b>State management</b> | Manual buffers       | Automatic <code>state_dict()</code> |

| Feature           | Your Implementation | PyTorch                                       |
|-------------------|---------------------|-----------------------------------------------|
| <b>Optimizers</b> | SGD, Adam, AdamW    | 10+ algorithms<br>(RMSprop, Adagrad,<br>etc.) |

### 11.6.2 Code Comparison

Optimizer usage in TinyTorch and PyTorch is nearly identical — by design. The patterns you learn here are the patterns you use in production.

## 11.7 Your TinyTorch

```
from tinytorch.core.optimizers import Adam

# Create optimizer for model parameters
optimizer = Adam(model.parameters(), lr=0.001)

# Training step
loss = criterion(predictions, targets)
loss.backward() # Compute gradients
optimizer.step() # Update parameters
optimizer.zero_grad() # Clear gradients
```

## 11.8 PyTorch

```
import torch.optim as optim

# Create optimizer for model parameters
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training step
loss = criterion(predictions, targets)
loss.backward() # Compute gradients
optimizer.step() # Update parameters
optimizer.zero_grad() # Clear gradients
```

Let's walk through each line to understand the comparison:

- **Line 1 (Import):** TinyTorch exposes optimizers from `tinytorch.core.optimizers`; PyTorch uses `torch.optim`. The namespace structure mirrors production frameworks.
- **Line 4 (Creation):** Both use identical syntax: `Adam(model.parameters(), lr=0.001)`. The `model.parameters()` method returns an iterable of tensors with `requires_grad=True`.
- **Line 7-8 (Training):** The loss computation and backward pass are identical. Your autograd system from Module 06 computes gradients just like PyTorch.
- **Line 9 (Update):** Both call `optimizer.step()` to update parameters using computed gradients. The update rules are mathematically identical.

- **Line 10 (Clear):** Both call `optimizer.zero_grad()` to clear gradients before the next iteration. Without this, gradients would accumulate across batches.

#### 💡 What's Identical

The optimizer API, update algorithms, and memory patterns are identical. When you debug Adam's learning rate or analyze optimizer memory usage in production, you'll understand exactly what's happening because you built these mechanisms yourself.

### 11.8.1 Why Optimizers Matter at Scale

Optimizer choice is a memory decision before it is a math decision. Three concrete data points:

- **Large language models (175B parameters):** Adam's optimizer state alone consumes **1.27 TB** (3x the 651.9 GB of parameters), forcing multi-GPU state sharding (e.g., DeepSpeed ZeRO).
- **Transformer training:** AdamW with `weight_decay=0.01` is the default, typically improving generalization over plain Adam by 2-5% accuracy.
- **Convergence speed:** Adam reaches target loss in roughly 30-50% fewer steps than SGD on vision and language tasks, paying back its higher per-step cost in wall-clock time.

When a model barely fits in memory with SGD, switching to Adam (1.5x more state) can be the line between “trains on one GPU” and “needs sharding.”

## 11.9 Check Your Understanding

### 💡 Check Your Understanding — Optimizers

Before moving on, verify you can articulate each of the following:

- Why Adam carries  $2\times$  the parameter memory compared to plain SGD (separate  $m$  and  $v$  buffers) and when that trade-off pays off in wall-clock time.
- Why the optimizer step is memory-bound on modern GPUs — streaming  $m$ ,  $v$ , params, and grads through VRAM saturates bandwidth long before compute is the bottleneck.
- What bias correction actually fixes (zero-initialized moments in early steps) and why it silently vanishes after a few hundred iterations.
- Why AdamW decouples weight decay from the adaptive scaling step — and why applying  $\lambda$  inside the gradient (as plain Adam does) under-regularizes parameters with large gradients.

If any of these feels fuzzy, revisit the Adam and AdamW sections before moving on.

Five questions. Try each one yourself before opening the answer — the math is small but the intuitions are the ones you will use in production.

#### Q1: Memory Calculation

A language model has 10 billion float32 parameters. Using Adam optimizer, how much total memory does optimizer state require? How does this compare to SGD with momentum?

### 💡 Answer

**Parameters:**  $10B \times 4 \text{ bytes} = 37.25 \text{ GB}$

**Adam state:** 2 buffers ( $m$ ,  $v$ ) =  $2 \times 37.25 \text{ GB} = 74.51 \text{ GB}$  **Total with Adam:**  $37.25 \text{ GB (params)} + 74.51 \text{ GB (state)} = 111.76 \text{ GB}$



**SGD with momentum:** 1 buffer (velocity) = **37.25 GB Total with SGD:** 37.25 GB (params) + 37.25 GB (state) = **74.51 GB**

**Difference:** Adam uses **37.25 GB more** than SGD (50% increase). On 80 GB H100s that is the difference between fitting on one GPU and needing two — or implementing optimizer state sharding.

### Q2: Convergence Trade-off

If Adam converges in 100,000 steps and SGD needs 200,000 steps, but Adam's per-step time is 1.2x slower due to additional computations, which optimizer finishes training faster?

#### 💡 Answer

**Adam:** 100,000 steps  $\times$  1.2 = **120,000 time units** **SGD:** 200,000 steps  $\times$  1.0 = **200,000 time units**

**Adam finishes 1.67x faster** despite the slower per-step cost. The 2x reduction in steps outweighs the 1.2x per-step overhead.

This is why Adam wins in practice: wall-clock time to convergence is what users pay for, not per-step efficiency.

### Q3: Bias Correction Impact

In Adam, bias correction divides the first moment by  $(1 - \beta^t)$ . At step 1 with  $\beta=0.9$ , the correction factor is 0.1. At step 10, it is 0.651. How does this affect early vs late training?

#### 💡 Answer

**Step 1:** divide by 0.1 = multiply by **10x** (huge correction) **Step 10:** divide by 0.651 = multiply by **1.54x** (moderate correction) **Step 100:** divide by 0.99997  $\approx$  multiply by **1.0x** (negligible correction)

**Early training:** Large corrections inflate the still-tiny moment estimates up to a sensible magnitude, so the optimizer can take meaningful steps from the very first iteration.

**Late training:** Corrections converge to 1.0 and silently disappear; Adam runs on raw moments.

**Without correction:**  $m$  starts at zero, so the first updates are roughly 10x smaller than intended — training crawls until the running averages catch up.

### Q4: Weight Decay Comparison

Adam adds weight decay to gradients before adaptive scaling. AdamW applies it after. For a parameter with  $\text{grad}=0.001$  and  $\text{param}=1.0$ , which experiences stronger regularization with  $\text{weight\_decay}=0.01$  and  $\text{lr}=0.1$ ?

#### 💡 Answer

**Adam approach:** - Modified  $\text{grad} = 0.001 + 0.01 \times 1.0 = 0.011$  - This gradient gets adaptively scaled (divided by  $\sqrt{v}$ , which is small for small gradients) - Effective decay is amplified by adaptive scaling

**AdamW approach:** - Pure gradient update uses  $\text{grad}=0.001$  (small adaptive step) - Then  $\text{param} = \text{param} \times (1 - 0.1 \times 0.01) = \text{param} \times 0.999$  (fixed 0.1% shrinkage)

**AdamW has consistent 0.1% weight decay** regardless of gradient magnitude. Adam's decay strength varies with adaptive learning rate scaling, making it inconsistent across parameters. AdamW's consistency leads to better regularization behavior.

### Q5: Optimizer State Checkpointing

You're training with Adam and checkpoint every 1000 steps. The checkpoint saves parameters and optimizer state ( $m$ ,  $v$  buffers). If you resume from step 5000 but change learning rate from 0.001 to 0.0001, should you restore old optimizer state or reset it?

 **Answer**

**Restore state (recommended):** The `m` and `v` buffers contain valuable information about gradient statistics accumulated over 5000 steps. Resetting loses this and causes the optimizer to “forget” learned gradient scales.

**Impact of restoring:** - Keeps adaptive learning rates calibrated to parameter-specific gradient magnitudes - Prevents slow re-convergence that happens when resetting - Learning rate change affects step size but not the adaptive scaling

**When to reset:** - If switching optimizer types (SGD → Adam) - If gradient distribution has fundamentally changed (switching datasets) - If debugging and suspecting corrupted state

**Production practice:** Always restore optimizer state when resuming training unless you have specific reasons to reset. The state is part of what makes Adam effective.

## 11.10 Key Takeaways

- **Optimizer choice is a memory decision:** SGD adds  $1\times$  parameter memory (velocity); Adam adds  $2\times$  (moments). For a 175B-parameter model, that is the difference between 652 GB and 1.27 TB of optimizer state alone.
- **The step is memory-bound, not compute-bound:** Adam’s update is element-wise arithmetic over massive buffers; VRAM bandwidth, not FLOPs, sets the floor on step latency.
- **Bias correction matters only in early training:** The  $1 - \beta^t$  term inflates still-tiny moments so Adam takes meaningful steps from iteration 1; it converges to 1.0 and disappears after a few hundred steps.
- **AdamW is the default for a reason:** Decoupling weight decay from the adaptive divisor gives every parameter the same regularization strength, independent of gradient magnitude — which is why every modern transformer trains with it.

**Coming next:** An optimizer that takes one step is not a training system. Module 08 wraps the inner loop in epochs, schedules, clipping, evaluation, and checkpointing — the orchestration that turns one `step()` into a million without human babysitting.

## 11.11 Further Reading

For students who want to understand the academic foundations and mathematical underpinnings of optimization algorithms:

### 11.11.1 Seminal Papers

- **Adam: A Method for Stochastic Optimization** - Kingma & Ba (2015). The original Adam paper introducing adaptive moment estimation with bias correction. Explains the motivation and derivation.
  - **Systems Implication:** Introduced the dual-moment state buffers that profoundly shifted the memory footprint of deep learning. By trading VRAM capacity for convergence speed, Adam forced systems engineers to invent techniques like ZeRO (Zero Redundancy Optimizer) to shard optimizer states across GPU clusters. [arXiv:1412.6980](https://arxiv.org/abs/1412.6980)
- **Decoupled Weight Decay Regularization (AdamW)** - Loshchilov & Hutter (2019). Identifies the weight decay bug in Adam and proposes the decoupled fix. Shows significant improvements on image classification and language modeling.
  - **Systems Implication:** Demonstrated that the mathematical formulation of regularization heavily impacts parameter update dynamics. Decoupling weight decay standardized the L2 penalty

application, cementing AdamW as the default optimizer for modern transformer training architectures. [arXiv:1711.05101](#)

- **On the Importance of Initialization and Momentum in Deep Learning** - Sutskever et al. (2013). Classic paper explaining why momentum works and how it accelerates convergence in deep networks.
  - **Systems Implication:** Established that retaining historical gradient trajectory (velocity) is essential for escaping pathological curvature. This introduced the first major optimizer state buffer, setting the precedent for stateful gradient descent algorithms. [ICML 2013](#)

### 11.11.2 Additional Resources

- **Tutorial:** [“An overview of gradient descent optimization algorithms”](#) by Sebastian Ruder - Comprehensive survey covering SGD variants, momentum methods, and adaptive learning rate algorithms
- **Documentation:** [PyTorch Optimization Documentation](#) - See how production frameworks organize and document optimization algorithms

## 11.12 What's Next

### Coming Up: Module 08 - Training

You have an optimizer that takes one step. Module 08 wraps it in the loop that takes a million: epochs, validation, checkpointing, learning rate schedules, and early stopping. The question it answers is the one this chapter raised but did not solve — *how do you actually drive a model to convergence without babysitting it?*

#### Preview - How Your Optimizers Get Used in Future Modules:

Table 11.5 traces how this module is reused by later parts of the curriculum.

Table 11.5: How optimizers feed into subsequent training modules.

| Module                  | What It Does            | Your Optimizers In Action                                              |
|-------------------------|-------------------------|------------------------------------------------------------------------|
| <b>08: Training</b>     | Complete training loops | <code>for epoch in range(10): loss.backward(); optimizer.step()</code> |
| <b>09: Convolutions</b> | Convolutional networks  | AdamW optimizes millions of CNN parameters efficiently                 |
| <b>13: Transformers</b> | Attention mechanisms    | Large models require careful optimizer selection                       |

## 11.13 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

 **Save Your Progress**

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.

## 🔥 Chapter 12

# Module 08: Training

The inner loop is four lines: forward, loss, backward, step. A training system is everything wrapped around it. Schedules, clipping, evaluation modes, and checkpoints are what turn one `step()` into a million without human babysitting. This is where the framework becomes something you can leave running overnight.

### Module Info

**FOUNDATION TIER** | Difficulty: ●●○○ | Time: 5-7 hours | Prerequisites: 01-07

This is the capstone of the Foundation Tier. The seven components you built — tensors, activations, layers, losses, dataloader, autograd, optimizers — finally fit together into a Trainer that actually learns.

## 12.1 Overview

Seven modules of components, none of them yet doing anything together. A tensor on its own does not learn. A layer on its own does not learn. Even autograd plus an optimizer does not learn — somebody has to call them, in order, on real data, again and again. That somebody is the training loop, and you are about to build it.

The loop itself is four lines: forward pass, loss, backward pass, optimizer step. Repeat. The hard part is what production wraps around those four lines. Learning rates need to start high and decay. Gradients sometimes explode and need clipping. Long runs crash and need checkpoints to resume from. Models need separate train and evaluation modes so dropout and batch norm behave correctly. By the end of this module you will have a `Trainer` class that handles all of this — the same architecture PyTorch Lightning and Hugging Face Transformers expose to millions of users, just smaller and yours.

## 12.2 Learning Objectives

### By completing this module, you will:

- **Implement** a complete Trainer class orchestrating forward pass, loss computation, backward pass, and parameter updates
- **Master** learning rate scheduling with cosine annealing that adapts training speed over time
- **Understand** gradient clipping by global norm that prevents training instability
- **Build** checkpointing systems that save and restore complete training state for fault tolerance
- **Analyze** training memory overhead ( $4-6\times$  model size) and checkpoint storage costs

## 12.3 What You'll Build

Five pieces wrapped around the inner loop: a learning-rate schedule, a gradient clipper, the training loop itself, an evaluation pass, and checkpoint save/load.

**Implementation roadmap:**




Figure 12.1: **TinyTorch Training Ecosystem:** Orchestration of scheduling, clipping, and checkpointing within the Trainer class.

Table 12.1 lays out the implementation in order, one part at a time.

Table 12.1: **Implementation roadmap for the Trainer and learning-rate schedule.**

| Part | What You'll Implement                       | Key Concept                              |
|------|---------------------------------------------|------------------------------------------|
| 1    | <code>CosineSchedule</code> class           | Learning rate annealing (fast → slow)    |
| 2    | <code>clip_grad_norm()</code> function      | Global gradient clipping for stability   |
| 3    | <code>Trainer.train_epoch()</code>          | Complete training loop with scheduling   |
| 4    | <code>Trainer.evaluate()</code>             | Evaluation mode without gradient updates |
| 5    | <code>Trainer.save/load_checkpoint()</code> | Training state persistence               |

The pattern you'll enable:

```
# Complete training pipeline (modules 01-07 working together)
trainer = Trainer(model, optimizer, loss_fn, scheduler, grad_clip_norm=1.0)
for epoch in range(100):
    train_loss = trainer.train_epoch(train_data)
    eval_loss, accuracy = trainer.evaluate(val_data)
    trainer.save_checkpoint(f"checkpoint_{epoch}.pkl")
```

### 12.3.1 What You're NOT Building (Yet)

To keep the module focused, you will **not** implement:

- Distributed training across multiple GPUs (PyTorch uses `DistributedDataParallel`)

- Mixed-precision training (PyTorch’s Automatic Mixed Precision relies on dedicated FP16/BF16 tensor types)
- Exotic schedulers — warmup, cyclic, one-cycle, polynomial decay (production frameworks ship dozens)

**You are building the core training orchestration.** That orchestration is what the rest of the framework plugs into.

## 12.4 API Reference

The signatures you need to satisfy. Keep this open in a second tab while you implement.

### 12.4.1 CosineSchedule

```
CosineSchedule(max_lr=0.1, min_lr=0.01, total_epochs=100)
```

Cosine annealing learning rate schedule that smoothly decreases from `max_lr` to `min_lr` over `total_epochs`.

Table 12.2 lists the schedule API.

Table 12.2: Method on the `CosineSchedule` class.

| Method              | Signature                                   | Description                           |
|---------------------|---------------------------------------------|---------------------------------------|
| <code>get_lr</code> | <code>get_lr(epoch: int) -&gt; float</code> | Returns learning rate for given epoch |

### 12.4.2 Gradient Clipping

```
clip_grad_norm(parameters: List, max_norm: float = 1.0) -> float
```

Clips gradients by global norm to prevent exploding gradients. Returns original norm for monitoring.

### 12.4.3 Trainer

```
Trainer(model, optimizer, loss_fn, scheduler=None, grad_clip_norm=None)
```

Orchestrates complete training lifecycle with forward pass, loss computation, backward pass, optimization, and checkpointing.

#### Core Methods

Table 12.3 lists the core methods on the trainer.

Table 12.3: Core methods on the `Trainer` class.

| Method                   | Signature                                                              | Description                               |
|--------------------------|------------------------------------------------------------------------|-------------------------------------------|
| <code>train_epoch</code> | <code>train_epoch(dataloader, accumulation_steps=1) -&gt; float</code> | Train for one epoch, returns average loss |
| <code>evaluate</code>    | <code>evaluate(dataloader) -&gt; Tuple[float, float]</code>            | Evaluate model, returns (loss, accuracy)  |

| Method          | Signature                          | Description                      |
|-----------------|------------------------------------|----------------------------------|
| save_checkpoint | save_checkpoint(path: str) -> None | Save complete training state     |
| load_checkpoint | load_checkpoint(path: str) -> None | Restore training state from file |

## 12.5 Core Concepts

Five ideas. The training loop, epochs vs. iterations, train vs. eval mode, learning rate scheduling, and gradient clipping. Every ML framework — yours, PyTorch, JAX — implements these the same way at the conceptual level.

### 12.5.1 The Training Loop

The training loop is a four-step pattern repeated thousands of times: push data through the model (forward pass), measure how wrong it is (loss), compute how to improve each parameter (backward pass), apply the update (optimizer step). Random weights at iteration 0; a network that recognizes digits at iteration 30,000.

Here's the complete training loop from your Trainer implementation:

The code in ?@lst-08-training-train-epoch makes this concrete.

```
def train_epoch(self, dataloader, accumulation_steps=1):
    """Train for one epoch through the dataset."""
    self.model.training = True
    self.training_mode = True

    total_loss = 0.0
    num_batches = 0
    accumulated_loss = 0.0

    for batch_idx, (inputs, targets) in enumerate(dataloader):
        # Forward pass
        outputs = self.model.forward(inputs)
        loss = self.loss_fn.forward(outputs, targets)

        # Scale loss for accumulation
        scaled_loss = loss.data / accumulation_steps
        accumulated_loss += scaled_loss

        # Backward pass
        loss.backward()

        # Update parameters every accumulation_steps
        if (batch_idx + 1) % accumulation_steps == 0:
            # Gradient clipping
            if self.grad_clip_norm is not None:
                params = self.model.parameters()
                clip_grad_norm(params, self.grad_clip_norm)
```



```

        # Optimizer step
        self.optimizer.step()
        self.optimizer.zero_grad()

        total_loss += accumulated_loss
        accumulated_loss = 0.0
        num_batches += 1
        self.step += 1

    # Handle remaining accumulated gradients
    if accumulated_loss > 0:
        if self.grad_clip_norm is not None:
            params = self.model.parameters()
            clip_grad_norm(params, self.grad_clip_norm)

        self.optimizer.step()
        self.optimizer.zero_grad()
        total_loss += accumulated_loss
        num_batches += 1

    avg_loss = total_loss / max(num_batches, 1)
    self.history['train_loss'].append(avg_loss)

    # Update scheduler
    if self.scheduler is not None:
        current_lr = self.scheduler.get_lr(self.epoch)
        self.optimizer.lr = current_lr
        self.history['learning_rates'].append(current_lr)

    self.epoch += 1
    return avg_loss

```

: Listing 8.1 — `Trainer.train_epoch()` orchestrates the inner loop: forward, loss, backward, clip, step, zero-grad, with gradient accumulation and scheduler update. {#lst-08-training-train-epoch}

Each iteration processes one batch: the model turns inputs into predictions, the loss function scores them against targets, the backward pass fills `.grad` on every parameter, the clipper bounds the global norm, and the optimizer applies the step. Total loss divided by batch count is the average training loss you watch to see convergence.

While this implementation reads like standard sequential Python, scaling this loop to production hardware transforms it into a highly asynchronous, choreographed dance between the host CPU (dispatching kernel instructions) and the target GPU (executing them massively in parallel).

#### **i** Systems Implication: CPU/GPU Sync Overhead

Innocuous operations like calculating `total_loss += loss.item()` or printing loss values inside the inner loop force a hard synchronization block. The CPU must halt and wait for the GPU to drain its computation queue just to copy a single scalar value back across the PCIe bus. This catastrophic synchronization stalls the entire training pipeline, starving the GPU. High-performance systems metic-

ulously defer, asynchronously accumulate, or strictly periodically sample these metrics to prevent interrupting the compute symphony.

To circumvent hard physical VRAM limits without sacrificing statistical stability, the `accumulation_steps` parameter introduces a profound memory-compute trade-off. If an architecture requires an effective batch size of 128 but the GPU can only harbor 32 samples at a time, setting `accumulation_steps=4` serially accrues gradients over four distinct forward-backward passes. This achieves the mathematically identical optimizer update of a 128-sample batch, trading raw wall-clock time for an artificially expanded memory horizon.

### 12.5.2 Epochs and Iterations

Training operates on two timescales: iterations (single batch updates) and epochs (complete passes through the dataset). The hierarchy is what lets you reason about training progress and cost.

An iteration processes one batch: forward, backward, step. With 10,000 samples and batch size 32, one epoch is 313 iterations ( $10,000 \div 32$ , rounded up). Convergence typically takes dozens to hundreds of epochs, so tens of thousands of iterations.

Scale that up and the implications get real. ImageNet has 1.2M images; batch size 256 and 90 epochs is 421,920 iterations ( $1,200,000 \div 256 \times 90$ ). At 250ms per iteration, that's **29 hours** on one GPU. The same arithmetic tells you whether a hyperparameter sweep is feasible or whether you should rent more machines.

Your Trainer tracks both: `self.step` counts total iterations across all epochs, while `self.epoch` counts how many complete dataset passes you've completed. Schedulers typically operate on epoch boundaries (learning rate changes each epoch), while monitoring systems track loss per iteration.

### 12.5.3 Train vs Eval Modes

Neural networks behave differently during training versus evaluation. Layers like dropout randomly zero activations during training (for regularization) but keep all activations during evaluation. Batch normalization computes running statistics during training but uses fixed statistics during evaluation. Your Trainer needs to signal which mode the model is in.

The pattern is simple: set `model.training = True` before training, set `model.training = False` before evaluation. This boolean flag propagates through layers, changing their behavior:

The code in `?@lst-08-training-evaluate` makes this concrete.

```
def evaluate(self, dataloader):
    """Evaluate model without updating parameters."""
    self.model.training = False
    self.training_mode = False

    total_loss = 0.0
    correct = 0
    total = 0

    for inputs, targets in dataloader:
        # Forward pass only (no backward!)
        outputs = self.model.forward(inputs)
        loss = self.loss_fn.forward(outputs, targets)

        total_loss += loss.data
```

```

# Calculate accuracy (for classification)
if len(outputs.data.shape) > 1: # Multi-class
    predictions = np.argmax(outputs.data, axis=1)
    if len(targets.data.shape) == 1: # Integer targets
        correct += np.sum(predictions == targets.data)
    else: # One-hot targets
        correct += np.sum(predictions == np.argmax(targets.data, axis=1))
    total += len(predictions)

avg_loss = total_loss / len(dataloader) if len(dataloader) > 0 else 0.0
accuracy = correct / total if total > 0 else 0.0

self.history['eval_loss'].append(avg_loss)

return avg_loss, accuracy

```

: Listing 8.2 — `Trainer.evaluate()` flips the model to eval mode, runs forward-only passes, and reports average loss and accuracy without touching parameters. {#lst-08-training-evaluate}

Notice what’s missing: no `loss.backward()`, no `optimizer.step()`, no gradient updates. Evaluation measures current model performance without changing parameters. This separation is crucial: if you accidentally left `training = True` during evaluation, dropout would randomly zero activations, giving you noisy accuracy measurements that don’t reflect true model quality.

### 12.5.4 Learning Rate Scheduling

Learning rate scheduling adapts training speed over time. Early in training, when parameters are far from optimal, high learning rates enable rapid progress. Late in training, when approaching a good solution, low learning rates enable stable convergence without overshooting. Fixed learning rates force you to choose between fast early progress and stable late convergence. Scheduling gives you both.

Cosine annealing uses the cosine function to smoothly transition from maximum to minimum learning rate:

```

def get_lr(self, epoch: int) -> float:
    """Get learning rate for current epoch."""
    if epoch >= self.total_epochs:
        return self.min_lr

    # Cosine annealing formula
    cosine_factor = (1 + np.cos(np.pi * epoch / self.total_epochs)) / 2
    return self.min_lr + (self.max_lr - self.min_lr) * cosine_factor

```

At epoch 0,  $\cos(0) = 1$  so `cosine_factor = 1.0` and the rate is `max_lr`. At the final epoch,  $\cos(\pi) = -1$  so `cosine_factor = 0.0` and the rate is `min_lr`. Between those endpoints the curve falls off smoothly — fast at first, slower as it bottoms out.

For `max_lr=0.1, min_lr=0.01, total_epochs=100`:

```

Epoch 0: 0.100    (aggressive learning)
Epoch 25: 0.087    (still fast)

```

```
Epoch 50: 0.055 (slowing down)
Epoch 75: 0.023 (fine-tuning)
Epoch 100: 0.010 (stable convergence)
```

Your Trainer applies the schedule automatically after each epoch:

```
if self.scheduler is not None:
    current_lr = self.scheduler.get_lr(self.epoch)
    self.optimizer.lr = current_lr
```

This updates the optimizer’s learning rate before the next epoch begins, creating adaptive training speed without manual intervention.

### 12.5.5 Gradient Clipping

Gradient clipping prevents exploding gradients that destroy training progress. During backpropagation, gradients sometimes become extremely large (thousands or even infinity), causing parameter updates that jump far from the optimum or overflow into NaN. Clipping rescales large gradients to a safe maximum while preserving their direction.

The key insight is clipping by global norm rather than individual gradients. Computing the norm across all parameters  $\sqrt{(\sum g^2)}$  and scaling uniformly preserves the relative magnitudes between different parameters:

The code in `?@lst-08-training-clip-grad-norm` makes this concrete.

```
def clip_grad_norm(parameters: List, max_norm: float = 1.0) -> float:
    """Clip gradients by global norm to prevent exploding gradients."""
    # Compute global norm across all parameters
    total_norm = 0.0
    for param in parameters:
        if param.grad is not None:
            grad_data = param.grad if isinstance(param.grad, np.ndarray) else
            param.grad.data
            total_norm += np.sum(grad_data ** 2)

    total_norm = np.sqrt(total_norm)

    # Scale all gradients if norm exceeds threshold
    if total_norm > max_norm:
        clip_coef = max_norm / total_norm
        for param in parameters:
            if param.grad is not None:
                if isinstance(param.grad, np.ndarray):
                    param.grad = param.grad * clip_coef
                else:
                    param.grad.data = param.grad.data * clip_coef

    return float(total_norm)
```

: Listing 8.3 — `clip_grad_norm()` rescales gradients uniformly when the global norm exceeds `max_norm`, preserving their relative magnitudes. {#lst-08-training-clip-grad-norm}

Consider gradients  $[100, 200, 50]$  with global norm  $\sqrt{(100^2 + 200^2 + 50^2)} \approx 229$ . With `max_norm=1.0`, the clip coefficient is  $1.0 / 229 \approx 0.00437$ , and every gradient is scaled by it:  $[0.437, 0.873, 0.218]$ . The new norm is exactly 1.0, but the relative magnitudes survive — the second gradient is still twice the first.

This uniform scaling is crucial. If we clipped each gradient independently to 1.0, we'd get  $[1.0, 1.0, 1.0]$ , destroying the information that the second parameter needs larger updates than the first. Global norm clipping prevents explosions while respecting the gradient's message about relative importance.

### 12.5.6 Checkpointing

Checkpointing saves complete training state to disk, enabling fault tolerance and experimentation. Training runs take hours or days. Hardware fails. You want to try different hyperparameters after epoch 50. Checkpoints make all of this possible by capturing everything needed to resume training exactly where you left off.

A complete checkpoint includes:

The code in `?lst-08-training-save-checkpoint` makes this concrete.

```
def save_checkpoint(self, path: str):
    """Save complete training state for resumption."""
    checkpoint = {
        'epoch': self.epoch,
        'step': self.step,
        'model_state': self._get_model_state(),
        'optimizer_state': self._get_optimizer_state(),
        'scheduler_state': self._get_scheduler_state(),
        'history': self.history,
        'training_mode': self.training_mode
    }

    Path(path).parent.mkdir(parents=True, exist_ok=True)
    with open(path, 'wb') as f:
        pickle.dump(checkpoint, f)
```

: Listing 8.4 — `Trainer.save_checkpoint()` pickles the full training state: parameters, optimizer buffers, scheduler, epoch/step counters, and history. {#lst-08-training-save-checkpoint}

Model state is straightforward: copy all parameter tensors. Optimizer state is more subtle: SGD with momentum stores velocity buffers (one per parameter), Adam stores two moment buffers (first and second moments). Scheduler state captures current learning rate progression. Training metadata includes epoch counter and loss history.

Loading reverses the process:

The code in `?lst-08-training-load-checkpoint` makes this concrete.

```
def load_checkpoint(self, path: str):
    """Restore training state from checkpoint."""
    with open(path, 'rb') as f:
        checkpoint = pickle.load(f)

    self.epoch = checkpoint['epoch']
    self.step = checkpoint['step']
```

```

self.history = checkpoint['history']
self.training_mode = checkpoint['training_mode']

# Restore states (simplified for educational purposes)
if 'model_state' in checkpoint:
    self._set_model_state(checkpoint['model_state'])
if 'optimizer_state' in checkpoint:
    self._set_optimizer_state(checkpoint['optimizer_state'])
if 'scheduler_state' in checkpoint:
    self._set_scheduler_state(checkpoint['scheduler_state'])

```

: Listing 8.5 — `Trainer.load_checkpoint()` reverses the save: unpickle, restore counters and history, then rehydrate model, optimizer, and scheduler state. {#lst-08-training-load-checkpoint}

After loading, training resumes as if the interruption never happened. The next `train_epoch()` call starts at the correct epoch, uses the correct learning rate, and continues optimizing from the exact parameter values where you stopped.

### 12.5.7 Computational Complexity

Training cost is a function of architecture and dataset size. For a fully connected network with  $L$  layers of width  $d$ , the forward pass is  $O(d^2 \times L)$  — matrix multiplications dominate. The backward pass has the same complexity, since autograd revisits every operation. With  $N$  samples and batch size  $B$ , one epoch is  $N / B$  iterations.

Total training cost for  $E$  epochs:

Time per iteration:  $O(d^2 \times L) \times 2$  (forward + backward)  
Iterations per epoch:  $N / B$   
Total iterations:  $(N / B) \times E$   
Total complexity:  $O((N \times E \times d^2 \times L) / B)$

Real numbers make this concrete. A 2-layer network ( $d=512$ ) on 10,000 samples (batch size 32) for 100 epochs:

|                    |                                |                                |
|--------------------|--------------------------------|--------------------------------|
| $d^2 \times L$     | $= 512^2 \times 2$             | $= 524,288$ ops per sample     |
| Batch operations   | $= 524,288 \times 32$          | $= 16.8\text{M}$ ops per batch |
| Iterations / epoch | $= 10,000 / 32$                | $= 313$                        |
| Total iterations   | $= 313 \times 100$             | $= 31,300$                     |
| Total operations   | $= 31,300 \times 16.8\text{M}$ | $\approx 525$ billion ops      |

At 1 GFLOP/s (typical CPU), that's about **525 seconds** ( $\approx 9$  minutes). A GPU at 1 TFLOP/s ( $1,000\times$  faster) finishes it in **0.5 seconds**. The arithmetic is exactly why GPUs exist for ML: the workload is dense linear algebra, and a GPU eats dense linear algebra.

Memory complexity is simpler but just as important:

Table 12.4 breaks down the memory footprint component by component.

Table 12.4: Memory footprint of model parameters, gradients, and optimizer state.

| Component        | Memory                                  |
|------------------|-----------------------------------------|
| Model parameters | $d^2 \times L \times 4$ bytes (float32) |
| Gradients        | Same as parameters                      |

| Component              | Memory                               |
|------------------------|--------------------------------------|
| Optimizer state (SGD)  | Same as parameters (momentum)        |
| Optimizer state (Adam) | 2× parameters (two moments)          |
| Activations            | $d \times B \times L \times 4$ bytes |

Total training memory is typically 4–6× model size, depending on optimizer. This explains GPU memory constraints: a 1GB model requires 4-6GB GPU memory for training, limiting batch size when memory is scarce.

## 12.6 Production Context

### 12.6.1 Your Implementation vs. PyTorch

Your `Trainer` and PyTorch’s training stack (Lightning, Hugging Face `Trainer`) share the same architecture. Production adds distributed training, mixed precision, dozens of schedulers, and a callback system. The inner loop is identical.

Table 12.5 places your implementation side by side with the production reference for direct comparison.

Table 12.5: Feature comparison between TinyTorch Trainer and PyTorch training stacks.

| Feature                     | Your Implementation          | PyTorch / Lightning                   |
|-----------------------------|------------------------------|---------------------------------------|
| <b>Training Loop</b>        | Manual forward/backward/step | Same pattern, with callbacks          |
| <b>Schedulers</b>           | Cosine annealing             | 20+ schedulers (warmup, cyclic, etc.) |
| <b>Gradient Clipping</b>    | Global norm clipping         | Same algorithm, GPU-optimized         |
| <b>Checkpointing</b>        | Pickle-based state saving    | Same concept, optimized formats       |
| <b>Distributed Training</b> | Single device                | Multi-GPU, multi-node                 |
| <b>Mixed Precision</b>      | FP32 only                    | Automatic FP16/BF16                   |

### 12.6.2 Code Comparison

Equivalent training pipelines side by side. Same conceptual flow: build model, attach optimizer and scheduler, hand them to a trainer, loop.

## 12.7 Your TinyTorch

```
from tinytorch import Trainer, CosineSchedule, SGD, MSELoss

# Setup
model = MyModel()
optimizer = SGD(model.parameters(), lr=0.1, momentum=0.9)
scheduler = CosineSchedule(max_lr=0.1, min_lr=0.01, total_epochs=100)
trainer = Trainer(model, optimizer, MSELoss(), scheduler, grad_clip_norm=1.0)
```

```
# Training loop
for epoch in range(100):
    train_loss = trainer.train_epoch(train_data)
    eval_loss, acc = trainer.evaluate(val_data)

    if epoch % 10 == 0:
        trainer.save_checkpoint(f"ckpt_{epoch}.pkl")
```

## 12.8 PyTorch Lightning

```
import torch
from torch.optim.lr_scheduler import CosineAnnealingLR
from pytorch_lightning import Trainer

# Setup (nearly identical!)
model = MyModel()
optimizer = torch.optim.SGD(model.parameters(), lr=0.1, momentum=0.9)
scheduler = CosineAnnealingLR(optimizer, T_max=100, eta_min=0.01)
trainer = Trainer(max_epochs=100, gradient_clip_val=1.0)

# Training (abstracted by Lightning)
trainer.fit(model, train_dataloader, val_dataloader)
# Lightning handles the loop, checkpointing, and callbacks automatically
```

The pieces line up almost one-for-one:

- **Imports** — TinyTorch exposes classes directly; PyTorch uses a deeper module hierarchy. Same concepts, different organization.
- **Model + optimizer** — identical pattern. Both pass `model.parameters()` to the optimizer so it can track what to update.
- **Scheduler** — `CosineSchedule` versus `CosineAnnealingLR`. Different names, same math.
- **Trainer setup** — TinyTorch takes model, optimizer, loss, and scheduler explicitly. Lightning hides them behind a model definition + `Trainer(...)`. Both support `gradient_clip`.
- **Training loop** — TinyTorch makes the epoch loop explicit; Lightning hides it inside `trainer.fit()`. The loop Lightning runs is the loop you wrote.
- **Checkpointing** — TinyTorch requires manual `save_checkpoint()` calls; Lightning checkpoints automatically based on validation metrics.

### 💡 What's Identical

The core training loop pattern: forward pass → loss → backward → gradient clipping → optimizer step → learning rate scheduling. When debugging PyTorch training, you'll understand exactly what's happening because you built it yourself.

### 12.8.1 Why Training Infrastructure Matters at Scale

The same patterns scale up brutally. Three numbers from production:



- **GPT-3 training:** 175 billion parameters, 300 billion tokens, ~\$4.6 million in compute. A single checkpoint is **350 GB** — larger than most laptop SSDs. Checkpoint frequency is a real engineering tradeoff between fault tolerance and storage cost.
- **ImageNet training:** 1.2 million images, 90 epochs is standard. At 250ms per iteration (batch 256), that's **29 hours** on one GPU. Learning rate scheduling is the difference between 75% accuracy (mediocre) and 76.5% (state-of-the-art) — a difference papers are written about.
- **Training instability:** Without gradient clipping, roughly 1 in 50 training runs randomly diverges — gradients explode, outputs go NaN, all progress lost. A 2% failure rate is unacceptable when each run costs thousands of dollars.

Your `Trainer` handles all three of these at educational scale, with the same architecture the production systems use. The numbers get bigger; the loop does not.

## 12.9 Check Your Understanding

### 💡 Check Your Understanding — Training

Before moving on, verify you can articulate each of the following:

- The CPU/GPU sync overhead introduced by a naive training loop — why `loss.item()` or a `print` statement inside the inner loop forces the GPU to drain its queue and stalls the whole pipeline.
- Why gradient accumulation trades wall-clock time for effective batch size without changing peak memory, and why that is exactly the knob you reach for when VRAM is the binding constraint.
- Why global-norm clipping preserves the *relative* magnitudes of gradients (which individual per-parameter clipping destroys) and why that distinction matters for learning.
- Why training memory is typically 4–6× the model size (params + grads + optimizer state + activations) and how that sets the floor on what fits on a single GPU.

If any of these feels fuzzy, revisit the Training Loop and Computational Complexity sections before moving on.

Five systems-thinking questions. They build intuition for the performance characteristics and trade-offs you'll meet again in production ML — usually under a deadline.

#### Q1: Training Memory Calculation

You have a model with 10 million parameters (float32) and use the Adam optimizer. Estimate total training memory required: parameters + gradients + optimizer state. Then compare with SGD.

### 💡 Answer

#### Adam optimizer:

- Parameters:  $10\text{M} \times 4 \text{ bytes} = \mathbf{40 \text{ MB}}$
- Gradients:  $10\text{M} \times 4 \text{ bytes} = \mathbf{40 \text{ MB}}$
- Adam state (two moments):  $10\text{M} \times 2 \times 4 \text{ bytes} = \mathbf{80 \text{ MB}}$
- **Total: 160 MB** (4× parameter size)

#### SGD with momentum:

- Parameters:  $10\text{M} \times 4 \text{ bytes} = \mathbf{40 \text{ MB}}$
- Gradients:  $10\text{M} \times 4 \text{ bytes} = \mathbf{40 \text{ MB}}$
- Momentum buffer:  $10\text{M} \times 4 \text{ bytes} = \mathbf{40 \text{ MB}}$
- **Total: 120 MB** (3× parameter size)

**Key insight:** Optimizer choice changes training memory by 33%. For large models close to the GPU's memory ceiling, SGD is sometimes the only optimizer that fits.

### Q2: Gradient Accumulation Trade-off

You want batch size 128 but your GPU can only fit 32 samples. You use gradient accumulation with `accumulation_steps=4`. How does this affect: (a) Memory usage? (b) Training time? (c) Gradient noise?

#### 💡 Answer

**(a) Memory:** No change. Only one batch (32 samples) in GPU memory at a time. Gradients accumulate in parameter `.grad` buffers which already exist.

**(b) Training time: 4× slower per update.** You process 4 batches sequentially (forward + backward) before optimizer step. Total iterations stays the same, but wall-clock time increases linearly with accumulation steps.

**(c) Gradient noise: Reduced** (same as true `batch_size=128`). Averaging gradients over 128 samples gives more accurate gradient estimate than 32 samples, leading to more stable training.

**Trade-off summary:** Gradient accumulation exchanges compute time for effective batch size when memory is limited. You get better gradients (less noise) but slower training (more time per update).

### Q3: Learning Rate Schedule Analysis

Training with fixed `lr=0.1` converges quickly initially but oscillates around the optimum, never quite reaching it. Training with cosine schedule (`0.1 → 0.01`) converges slower initially but reaches better final accuracy. Explain why, and suggest when fixed LR might be better.

#### 💡 Answer

**Why fixed LR oscillates:** High learning rate (0.1) enables large parameter updates. Early in training (far from optimum), large updates accelerate convergence. Near the optimum, large updates overshoot, causing oscillation: update jumps past the optimum, then jumps back, repeatedly.

**Why cosine schedule reaches better accuracy:** Starting high (0.1) provides fast early progress. Gradual decay (`0.1 → 0.01`) allows the model to take progressively smaller steps as it approaches the optimum. By the final epochs, `lr=0.01` enables fine-tuning without overshooting.

**When fixed LR is better:** - **Short training runs** (< 10 epochs): Scheduling overhead not worth it -

**Learning rate tuning:** Finding optimal LR is easier with fixed values - **Transfer learning:** When fine-tuning pre-trained models, fixed low LR (0.001) often works best

**Rule of thumb:** For training from scratch over 50+ epochs, scheduling almost always improves final accuracy by 1-3%.

### Q4: Checkpoint Storage Strategy

You're training for 100 epochs. Each checkpoint is 1 GB. Checkpointing every epoch costs 100 GB of storage. Checkpointing every 10 epochs risks losing 10 epochs of work if training crashes. Design a checkpointing strategy that balances fault tolerance and storage cost.

#### 💡 Answer

**Strategy: Keep last N + best + milestones**

1. **Keep last N=3 checkpoints** (rolling window): `epoch_98.pkl`, `epoch_99.pkl`, `epoch_100.pkl` (3 GB)
2. **Keep best checkpoint** (lowest validation loss): `best_epoch_72.pkl` (1 GB)

3. **Keep milestone checkpoints** (every 25 epochs): epoch\_25.pkl, epoch\_50.pkl, epoch\_75.pkl (3 GB)

**Total storage: 7 GB** (vs 100 GB for every epoch)

**Fault tolerance:**

- Last 3 checkpoints: Lose at most 1 epoch of work
- Best checkpoint: Can always restart from best validation performance
- Milestones: Can restart experiments from quarter-points

**Implementation:**

```
if epoch % 25 == 0: # Milestone
    save_checkpoint(f"milestone_epoch_{epoch}.pkl")
elif epoch >= last_3_start: # Last 3
    save_checkpoint(f"recent_epoch_{epoch}.pkl")
if is_best_validation: # Best
    save_checkpoint(f"best_epoch_{epoch}.pkl")
```

**Production systems** use this strategy plus cloud storage for off-site backup.

#### Q5: Global Norm Clipping Analysis

Two training runs: (A) clips each gradient individually to max 1.0, (B) clips by global norm with  $\text{max\_norm}=1.0$ . Both encounter gradients  $[50, 100, 5]$  with global norm  $\sqrt{(50^2 + 100^2 + 5^2)} \approx 112$ . What are the clipped gradients in each case? Which preserves gradient direction better?

#### 💡 Answer

**(A) Individual clipping** (clip each to max 1.0):

- Original:  $[50, 100, 5]$
- Clipped:  $[1.0, 1.0, 1.0]$
- **Result:** All parameters get equal updates — relative-importance information is destroyed.

**(B) Global norm clipping** (scale uniformly):

- Original:  $[50, 100, 5]$ , global norm  $\approx 112$
- Scale factor:  $1.0 / 112 \approx 0.0089$
- Clipped:  $[0.45, 0.89, 0.04]$
- New global norm: **1.0** (exactly  $\text{max\_norm}$ )
- **Result:** Relative magnitudes preserved — the second parameter still gets a  $2\times$  larger update than the first.

**Why (B) is better:** Gradients encode relative importance: parameter 2 *needs* larger updates than parameter 1. Global-norm clipping bounds the explosion while respecting that signal. Individual clipping flattens the signal, treating every parameter as equally important.

**Verification:**  $\sqrt{(0.45^2 + 0.89^2 + 0.04^2)} \approx 1.0$ .

## 12.10 Key Takeaways

- **The inner loop is four lines; the Trainer is everything around them:** Forward, loss, backward, step repeats forever. Production wraps it in schedules, clipping, evaluation modes, and checkpointing — those are what separate a toy loop from a training system.

- **Sync points kill throughput:** A `.item()` call or a print inside the inner loop drags a scalar back across PCIe and halts the GPU. High-performance loops defer metrics or sample them periodically.
- **Gradient accumulation decouples memory from effective batch size:** 32 samples in memory, serially accumulated  $4\times$  gives the mathematically identical update of a 128-batch — trading wall-clock time for an artificially expanded memory horizon.
- **Global-norm clipping preserves gradient *direction* while bounding magnitude:** Individual clipping flattens the signal; global-norm clipping bounds explosions without lying about which parameters need larger updates.
- **Checkpoints are the fault-tolerance boundary:** Saving params + optimizer state + scheduler + epoch makes a 29-hour run resumable. Every minute between checkpoints is a minute you are willing to lose.

**Coming next:** You just shipped the Foundation Tier. Module 09 opens the Architecture Tier with `Conv2d`, `MaxPool2d`, and the structural prior that made computer vision work — while the Trainer you just built keeps driving the loop unchanged.

## 12.11 Further Reading

For students who want to understand the academic foundations and advanced training techniques:

### 12.11.1 Seminal Papers

- **Cyclical Learning Rates for Training Neural Networks** - Smith (2017). Introduced cyclical learning rate schedules and the learning rate finder technique. Cosine annealing is a variant of these ideas.
  - **Systems Implication:** By architecting non-monotonic learning rate schedules that enable significantly faster convergence, this work directly reduced the total end-to-end wall-clock time and the exorbitant compute-hour costs required to train large-scale models on datacenter hardware. [arXiv:1506.01186](#)
- **On the Difficulty of Training Recurrent Neural Networks** - Pascanu et al. (2013). Analyzed the exploding and vanishing gradient problem, introducing gradient clipping as a robust stabilization solution. The global norm clipping you implemented originates directly from this research.
  - **Systems Implication:** While mathematically stabilizing, global gradient clipping requires an exhaustive reduction operation across *all* parameter gradients. In distributed training, this introduces a mandatory, global cross-GPU communication block (an `AllReduce` operation) that can severely bottleneck the pipeline right before parameter updates. [arXiv:1211.5063](#)
- **Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour** - Goyal et al. (2017). Demonstrated the mathematical scaling laws linking batch size and learning rate, formalizing linear warmup and distributed gradient accumulation techniques.
  - **Systems Implication:** Provided the rigorous empirical proof that scaling batch sizes linearly allows deep models to be efficiently distributed across massive GPU clusters via Data Parallelism. This high-batch-size regime effectively masks network communication latency behind large, dense matrix multiplications, achieving peak hardware utilization. [arXiv:1706.02677](#)

### 12.11.2 Additional Resources

- **PyTorch Lightning Documentation:** [Training Loop Documentation](#) - See how production frameworks implement the same training patterns you built
- **Weights & Biases Tutorial:** “Hyperparameter Tuning” - Excellent guide on learning rate scheduling and gradient clipping in practice

## 12.12 What's Next

You just shipped the Foundation Tier. Tensors, autograd, optimizers, dataloader, and a `Trainer` that ties them together — these are the load-bearing pieces of every modern ML framework, and you wrote all of them. The next tier is about *what gets put inside the model*. The training loop you built does not change.

Before moving on, the next three chapters give those Foundation pieces their first real workout. The **Foundation Milestones** — Rosenblatt's 1958 Perceptron, the 1969 XOR Crisis, and the 1986 MLP Revival — are runnable recreations of the experiments that shaped early neural-network history, moving from your forward-pass stack to the full autograd, optimizer, and `Trainer` pipeline you just finished. You watch your own code fail in the same way Minsky proved it had to, then break through with the same fix Rumelhart shipped. Then, on the far side, Module 09 opens the Architecture Tier.

### Coming Up: Foundation Milestones, then Module 09 — Convolutions

First: three Foundation Milestones start with your early layers, expose the XOR limitation, then use your full `Trainer` for TinyDigits MLP recognition — proof that the framework you built reproduces the history of the field. Then Module 09 opens the Architecture Tier with `Conv2d`, `MaxPool2d`, and `Flatten`: the layers that exploit spatial structure in images and make computer vision possible. Same `Trainer.train_epoch()` you wrote here will train the CNNs you build there, with no code changes. That's the payoff of separating orchestration from architecture.

#### Preview — how the `Trainer` you just built gets reused:

Table 12.6 traces how this module is reused by later parts of the curriculum.

Table 12.6: How the `Trainer` gets reused in the Architecture tier modules.

| Module                  | What It Adds                      | Your <code>Trainer</code> In Action                   |
|-------------------------|-----------------------------------|-------------------------------------------------------|
| <b>09: Convolutions</b> | Spatial layers for images         | Same <code>train_epoch()</code> trains CNNs unchanged |
| <b>Milestone: MLP</b>   | Solve XOR and train on TinyDigits | <code>Trainer</code> orchestrates the full pipeline   |
| <b>Milestone: CNN</b>   | Complete CIFAR-10 classification  | Vision models trained with your infrastructure        |

## 12.13 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



---

PART II

# Foundation Milestones

---





## 🔥 Chapter 13

# Historical Milestones

### About This Part

**Proof-of-Mastery Demonstrations** | 6 Milestones | Prerequisites: Vary by milestone  
Milestones are runnable recreations of historical ML breakthroughs that use **YOUR** TinyTorch implementations. Each one validates that the components you built across the modules can reproduce results that once made headlines.

## 13.1 Overview

You’ve spent the modules building a working ML framework — tensors, autograd, layers, optimizers, attention. The milestones answer the only question that matters: does it actually run the experiments that defined the field?

You’ll find out by rebuilding history. Each milestone reproduces a landmark result — Rosenblatt’s Perceptron, the XOR crisis, backpropagation, convolutional networks, transformers, MLPerf — using *your* code. When the Perceptron computes its first predictions, it’s running your tensor, layer, and activation stack. When attention processes a sequence, it’s running your multi-head attention on top of your transformer block. When a CNN recognizes TinyDigits, those are your convolutional layers extracting the features.

That makes these chapters proof — to yourself, and to anyone reading your repo — that the framework you built is the same kind of artifact the original papers shipped on.

## 13.2 The Journey

Table 13.1 traces the historical milestone timeline and the modules each one requires.

Table 13.1: **Historical milestone timeline and required modules for each.**

| Year | Milestone      | What You’ll Build                                 | Required Modules |
|------|----------------|---------------------------------------------------|------------------|
| 1958 | Perceptron     | First neural network forward pass                 | 01–03            |
| 1969 | XOR Crisis     | Experience the AI Winter trigger                  | 01–03            |
| 1986 | MLP Revival    | Backprop solves XOR + digit recognition           | 01–08            |
| 1998 | CNN Revolution | TinyDigits CNN, with optional CIFAR-10 scale-up   | 01–09            |
| 2017 | Transformers   | Multi-head attention on structured sequence tasks | 01–08, 11–13     |
| 2018 | MLPerf         | Production optimization pipeline                  | 01–08, 14–19     |

### 13.3 Why Milestones Transform Learning

**You'll feel the historical struggle.** When your single-layer perceptron hits 50% accuracy on XOR and refuses to budge — loss stuck at 0.69, epoch after epoch — you'll understand in your bones why Minsky's proof stalled neural-network research for a decade. The AI Winter wasn't abstract skepticism; it was researchers watching their perceptrons fail in exactly the way yours just did.

**You'll experience the breakthrough.** Then you add one hidden layer. Same data, same training loop. Suddenly: 100% accuracy. Loss collapses to zero. You didn't just read about how depth unlocks non-linear representations — you watched your two-layer network solve what your one-layer network couldn't. That's lived experience, not summary.

**You'll build something real.** By Milestone 04 you're done with toy demos. You're training a LeNet-style CNN on TinyDigits, extracting spatial features with your convolutional layers, and optionally scaling the same code path to CIFAR-10 — using a network you wrote line by line, on a framework you wrote module by module.

### 13.4 How to Use Milestones

```
tito module status

tito milestone run 01

cd milestones/01_1958_perceptron
python 01_rosenblatt_forward.py
```

Each `tinytorch/milestones/NN_yyyy_name/` folder contains:

- **README.md** — full historical context and instructions
- **Python scripts** — runnable demonstrations for each milestone part

### 13.5 Learning Philosophy

```
Module teaches:  HOW  to build the component
Milestone proves: WHAT you can build with it
```

Modules give you the parts. Milestones force the parts to do real work — the same work that, in each case, moved the field forward.

### 13.6 What's Next?

Start at the beginning. Open `milestones/01_1958_perceptron/`, run `01_rosenblatt_forward.py`, and watch a single-layer network — built on your `tensor`, `Linear`, and `Sigmoid` implementations — compute real predictions before any training infrastructure exists. From there the path is chronological: each milestone exposes the next constraint, then later milestones introduce the ideas that broke through it.

Build the future by understanding the past.

## 🔥 Chapter 14

# Milestone 01: The Perceptron (1958)

## i Milestone Info

**Foundation Milestone** | Difficulty: ●○○○ | Time: 15–30 min | Prerequisites: Modules 01–03

## 💡 What You'll Learn

- Why random weights produce random results
- How tensors, activations, and layers form a perceptron
- Why training infrastructure is needed before models can learn

## 14.1 Overview

You just finished the forward-pass foundation. Your `Tensor` (Module 01), activations (Module 02), and `Linear` layer (Module 03) are working. This milestone runs the simplest possible model those pieces can drive — and the one that started the field.

It's 1958. Computers fill entire rooms and can barely add numbers. Then Frank Rosenblatt makes an outrageous claim: he's built a machine that can learn from examples. Not through hand-written rules, but through adjustable weights.

The press goes wild. The Navy funds research expecting machines that will “walk, talk, see, write, reproduce itself and be conscious of its existence.” The New York Times runs the headline: “*New Navy Device Learns by Doing.*”

The optimism was premature. The insight wasn't. You're about to recreate the moment machine learning was born — with components YOU built yourself.

## 14.2 What You'll Build

A single-layer perceptron for binary classification that demonstrates the forward path before training enters the curriculum:

1. **The Architecture** — tensors flow through a `Linear` layer and sigmoid activation
2. **The Limitation** — random weights produce random predictions (~50% accuracy)

```
Input (features) --> Linear --> Sigmoid --> Output (0 or 1)
```

## 14.3 Prerequisites

Table 14.1 lists the modules you need to have completed before starting.

Table 14.1: Prerequisite modules for the Perceptron milestone.

| Module | Component   | What It Provides        |
|--------|-------------|-------------------------|
| 01     | Tensor      | YOUR data structure     |
| 02     | Activations | YOUR sigmoid activation |
| 03     | Layers      | YOUR Linear layer       |

## 14.4 Running the Milestone

Finish Modules 01–03 first. Check your progress:

```
tito module status
```

```
cd milestones/01_1958_perceptron

# See the forward-pass problem
python 01_rosenblatt_forward.py
# Expected: ~50% accuracy (random guessing)

# Or run it from the TinyTorch project root
tito milestone run 01
```

## 14.5 Expected Results

Table 14.2 records the accuracy and runtime you should expect to see.

Table 14.2: Expected accuracy for the Perceptron milestone scripts.

| Script            | Accuracy | What It Shows                    |
|-------------------|----------|----------------------------------|
| 01 (Forward Only) | ~50%     | Random weights = random guessing |

## 14.6 The Aha Moment: Learning IS the Intelligence

You'll run a single forward-only script using YOUR Linear layer and YOUR sigmoid. It should land near 50% accuracy because the weights are still random.

**What's missing?** Not the model. Not the data. The learning loop.

```
output = model(input)          # YOUR code computes
loss = loss_fn(output, target) # YOUR code measures
# No backward(), no optimization, no learning
# Result: Random weights stay random
```

Run the script and watch YOUR Linear layer make random guesses — near coin-flip accuracy. That failure is the point: the architecture is ready, but learning requires the modules that come next.

You just recreated the forward pass of Rosenblatt’s perceptron. The next milestones add the training machinery that turns this architecture into a learner.

## 14.7 Your Code Powers This

Table 14.3 names the TinyTorch components that power this milestone.

Table 14.3: TinyTorch components that power the Perceptron milestone.

| Component | Your Module | What It Does               |
|-----------|-------------|----------------------------|
| Tensor    | Module 01   | Stores inputs and weights  |
| Sigmoid   | Module 02   | YOUR activation function   |
| Linear    | Module 03   | YOUR fully-connected layer |

## 14.8 Historical Context

Rosenblatt didn’t just publish — he built. The Mark I Perceptron was custom hardware: a  $20 \times 20$  grid of photocells wired to motor-driven potentiometers that physically adjusted the weights. The 1958 paper established the two ideas under every modern model: trainable weights and error-driven learning. Eleven years later, Minsky and Papert’s *Perceptrons* (1969) proved single-layer networks couldn’t learn XOR. Funding collapsed. The first AI winter began.

## 14.9 Systems Insights

- **Memory:**  $O(n)$  parameters for  $n$  input features
- **Compute:**  $O(n)$  operations per sample
- **Limitation:** Can only solve linearly separable problems

## 14.10 What’s Next

Linear separability — the Perceptron’s hard ceiling — sparked the first AI winter. **Milestone 02** runs your network on XOR and watches that limit appear.

## 14.11 Further Reading

- **Original Paper:** Rosenblatt, F. (1958). “The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain”
- **Wikipedia:** [Perceptron](#)



## 🔥 Chapter 15

# Milestone 02: The XOR Crisis (1969)

## i Milestone Info

**Foundation Milestone** | Difficulty: ●●○○ | Time: 30–45 min | Prerequisites: Modules 01–08

## 💡 What You'll Learn

- Why single-layer networks have fundamental mathematical limits
- How hidden layers enable non-linear decision boundaries
- Why “deep” learning is called DEEP

## 15.1 Overview

It's 1969. Neural networks are the hottest thing in AI. Funding is pouring in. Then Marvin Minsky and Seymour Papert publish a 308-page mathematical proof that destroys everything: perceptrons **cannot solve XOR**. Not “struggle with” — CANNOT. Mathematically impossible.

Funding evaporates overnight. Research labs shut down. The field dies for 17 years — the infamous **AI Winter**.

You're about to live that crisis. You'll watch your own perceptron — built from your own modules — fail on four points despite a flawless training loop. Loss stuck at 0.69. Accuracy frozen at 50%. Epoch after epoch of futility. Then you'll add one hidden layer and watch the impossible collapse into the trivial.

## 15.2 What You'll Build

Two demonstrations of perceptron limitations and the multi-layer solution:

1. **The Crisis** — watch a perceptron fail on XOR despite training
2. **The Solution** — add a hidden layer and solve the “impossible” problem

```
Crisis:   Input --> Linear --> Output (FAILS)
Solution: Input --> Linear --> ReLU --> Linear --> Output (100%!)
```

## 15.3 The XOR Problem

| Inputs |    | Output            |
|--------|----|-------------------|
| x1     | x2 | XOR               |
| 0      | 0  | --> 0 (same)      |
| 0      | 1  | --> 1 (different) |

```
1  0  -->  1  (different)
1  1  -->  0  (same)
```

Plot those four points. The two zeros sit on one diagonal, the two ones on the other. No straight line separates them — and a single-layer perceptron can only draw straight lines. No amount of training fixes that. It's geometry, not optimization.

## 15.4 Prerequisites

Table 15.1 lists the modules you need to have completed before starting.

Table 15.1: **Prerequisite modules for the XOR milestone.**

| Module | Component   | What It Provides               |
|--------|-------------|--------------------------------|
| 01     | Tensor      | YOUR data structure            |
| 02     | Activations | YOUR sigmoid/ReLU              |
| 03     | Layers      | YOUR Linear layers             |
| 04     | Losses      | YOUR loss functions            |
| 05     | DataLoader  | YOUR data pipeline             |
| 06     | Autograd    | YOUR automatic differentiation |
| 07     | Optimizers  | YOUR SGD optimizer             |
| 08     | Training    | YOUR training loop             |

## 15.5 Running the Milestone

Before running, ensure you have completed Modules 01–08. You can check your progress:

```
tito module status
```

```
cd milestones/02_1969_xor
```

```
# Part 1: Experience the crisis
```

```
python 01_xor_crisis.py
```

```
# Expected: Loss stuck at ~0.69, accuracy ~50%
```

```
# Part 2: See the solution
```

```
python 02_xor_solved.py
```

```
# Expected: Loss --> 0.0, accuracy 100%
```

## 15.6 Expected Results

Table 15.2 records the accuracy and runtime you should expect to see.



Table 15.2: Expected loss and accuracy for the XOR milestone scripts.

| Script            | Layers | Loss           | Accuracy | What It Shows          |
|-------------------|--------|----------------|----------|------------------------|
| 01 (Single Layer) | 1      | ~0.69 (stuck!) | ~50%     | Cannot learn XOR       |
| 02 (Multi-Layer)  | 2      | -> 0.0         | 100%     | Hidden layers solve it |

## 15.7 The Aha Moment: Depth Changes Everything

The numbers in the table are the *aftermath*. Live, the experiment feels different.

Script 01 starts training. Loss: 0.69... 0.69... 0.69. Still 0.69. Why isn't it learning? Did you break something?

You check the code. Everything's correct. Your Linear layer works. Your autograd computes gradients. Your optimizer updates weights. But accuracy stays at 50%.

Then it lands: it's not broken. It's *impossible*. This is what Minsky proved. This is why funding died. Your code is slamming into the same mathematical wall that nearly ended AI research — every component working perfectly, all of it useless against XOR's geometry.

Then you run script 02. Add one hidden layer. Loss drops immediately: 0.5... 0.3... 0.1... 0.01... 0.0. Accuracy: 100%.

**Depth enables non-linear decision boundaries.** The hidden layer learns to bend the input space until XOR becomes linearly separable. A single layer can only draw straight lines. Stack two, and you can draw any shape you need.

Same code. Same training loop. Same four points. The impossible is now trivial — and you've earned the right to call this *deep* learning.

## 15.8 Your Code Powers This

Table 15.3 names the TinyTorch components that power this milestone.

Table 15.3: TinyTorch components that power the XOR milestone.

| Component     | Your Module | What It Does                     |
|---------------|-------------|----------------------------------|
| Tensor        | Module 01   | Stores inputs and weights        |
| ReLU          | Module 02   | YOUR activation for hidden layer |
| Linear        | Module 03   | YOUR fully-connected layers      |
| BCELoss       | Module 04   | YOUR loss computation            |
| DataLoader    | Module 05   | YOUR data pipeline               |
| backward()    | Module 06   | YOUR autograd engine             |
| SGD           | Module 07   | YOUR optimizer                   |
| Training loop | Module 08   | YOUR training orchestration      |

## 15.9 Systems Insights

- **Memory:**  $O(n^2)$  with hidden layers (vs  $O(n)$  for perceptron)
- **Compute:**  $O(n^2)$  operations
- **Breakthrough:** Hidden representations unlock non-linear problems

## 15.10 Historical Context

Minsky and Papert's proof was mathematically airtight — and read as a verdict on the whole research program. Multi-layer networks were known, but no one had a practical way to train them. That gap took 17 years to close: Rumelhart, Hinton, and Williams published backpropagation through hidden layers in 1986, and the field exhaled.

The lesson is uncomfortable. A correct theorem, applied to the wrong abstraction, set an entire field back nearly two decades.

## 15.11 What's Next

XOR is a toy: four points, two dimensions, a problem you can solve in your head. The real question is whether the same trick — stack a hidden layer, let it learn its own representation — survives contact with image data. **Milestone 03** points the same training stack at TinyDigits and finds out.

## 15.12 Further Reading

- **The Crisis:** Minsky, M., & Papert, S. (1969). "Perceptrons: An Introduction to Computational Geometry"
- **The Solution:** Rumelhart, Hinton, Williams (1986). "[Learning representations by back-propagating errors](#)"
- **Wikipedia:** [AI Winter](#)

## 🔥 Chapter 16

# Milestone 03: The MLP Revival (1986)

## i Milestone Info

**Foundation Milestone** | Difficulty: ●●○○ | Time: 1–2 hours (incl. training) | Prerequisites: Modules 01–08

## 💡 What You'll Learn

- How a multilayer network discovers its own features (edges, strokes) with no hand-coding
- Why representation learning replaced manual feature engineering
- That YOUR ~100 lines of TinyTorch can solve XOR and train on TinyDigits

## 16.1 Overview

**For 17 years, neural networks were dead.**

Minsky's XOR proof (Milestone 02) showed a single layer of perceptrons could not separate even four points on a plane. Funding evaporated. Researchers moved on. "Neural network" became a dirty word.

Then in 1986, **Rumelhart, Hinton, and Williams** published "Learning representations by back-propagating errors." Their argument was structural: stack two layers with a nonlinearity between them, train every weight with the chain rule, and the network discovers its own features. No hand-crafted rules. No domain experts. Data in, patterns out.

This milestone recreates that result in two release-friendly steps: first solve XOR with a hidden layer, then train on the shipped TinyDigits dataset.

## 16.2 What You'll Build

Multi-layer perceptrons (MLPs) for non-linear learning and digit recognition:

1. **XOR Solved** — hidden layers plus backpropagation solve the 1969 crisis
2. **TinyDigits** — quick proof-of-concept on 8×8 images

Images --> Flatten --> Linear --> ReLU --> Linear --> ReLU --> Linear --> Classes

## 16.3 Prerequisites

Table 16.1 lists the modules you need to have completed before starting.

Table 16.1: Prerequisite modules for the MLP milestone.

| Module | Component               | What It Provides                     |
|--------|-------------------------|--------------------------------------|
| 01–04  | Foundation              | Tensor, Activations, Layers, Losses  |
| 05     | DataLoader              | YOUR batching and data pipeline      |
| 06–08  | Training Infrastructure | Autograd, Optimizers, Training loops |

## 16.4 Running the Milestone

Confirm Modules 01–08 are complete:

```
tito module status
```

```
cd milestones/03_1986_mlp
```

```
# Quick validation on TinyDigits
python 01_rumelhart_tinydigits.py
# Expected: 85%+ accuracy
```

```
# Or run the complete milestone from the TinyTorch project root
tito milestone run 03
```

```
# Run individual parts
tito milestone run 03 --part 1 # XOR Solved
tito milestone run 03 --part 2 # TinyDigits
```

## 16.5 Expected Results

Table 16.2 records the accuracy and runtime you should expect to see.

Table 16.2: Expected accuracy and training time for the MLP milestone scripts.

| Script          | Dataset       | Parameters | Accuracy | Training Time |
|-----------------|---------------|------------|----------|---------------|
| 01 (XOR Solved) | 4 examples    | small      | 100%     | <1 min        |
| 02 (TinyDigits) | 1K train, 8×8 | ~2.4K      | 85%+     | 3–5 min       |

## 16.6 The Aha Moment: Automatic Feature Discovery

**Watch YOUR network learn something you never taught it.**

After training, reshape the first hidden layer’s weights into image-sized patches and visualize them. You will see edge detectors — horizontal, vertical, diagonal strokes. Nobody wrote those filters. The network discovered them because edges happen to be useful for telling digits apart.

This is **representation learning**: the model invents its own features from data instead of waiting for an expert to hand-design them. Combined with the universal approximation result — one hidden layer plus a

nonlinearity can approximate any continuous function — this is why a stack of `Linear` and `ReLU` layers can move from XOR to image classification.

Your ~100 lines of TinyTorch just replicated the breakthrough that ended the first AI winter.

## 16.7 Your Code Powers This

Every component comes from YOUR implementations:

Table 16.3 names the TinyTorch components that power this milestone.

Table 16.3: TinyTorch components that power the MLP milestone.

| Component                     | Your Module | What It Does                |
|-------------------------------|-------------|-----------------------------|
| <code>Tensor</code>           | Module 01   | Stores images and weights   |
| <code>Linear</code>           | Module 03   | YOUR fully-connected layers |
| <code>ReLU</code>             | Module 02   | YOUR activation functions   |
| <code>CrossEntropyLoss</code> | Module 04   | YOUR loss computation       |
| <code>DataLoader</code>       | Module 05   | YOUR batching pipeline      |
| <code>backward()</code>       | Module 06   | YOUR autograd engine        |
| <code>SGD</code>              | Module 07   | YOUR optimizer              |

No PyTorch. No TensorFlow. Just YOUR code learning to read handwritten digits.

## 16.8 Historical Context

Backpropagation made multi-layer networks practical after the XOR crisis. The same idea that solves four non-linear points also scales to handwritten digit recognition: every deep learning system you have used descends from that chain-rule machinery.

## 16.9 Systems Insights

- **Memory:**  $\sim 100\text{K parameters} \times 4 \text{ bytes} = 400 \text{ KB}$  of weights — small enough to fit on 1986 workstation RAM, which is partly why this experiment was even possible.
- **Compute:** Dense matrix multiplies dominate training time. Every forward pass through a fully-connected layer is one big GEMM.
- **Architecture:** Each hidden layer composes features from the layer below, building progressively more abstract representations as you stack depth.

## 16.10 What's Next

MLPs treat images as flat vectors. To your network, pixel (0,0) and pixel (0,1) are no more related than pixel (0,0) and pixel (27,27) — spatial structure is thrown away the moment you call `flatten`. **Milestone 04** (CNN) puts locality back in with convolutional layers, and asks what that costs and what it buys.

## 16.11 Further Reading

- **The Backprop Paper:** Rumelhart, Hinton, Williams (1986). “[Learning representations by back-propagating errors](#)”

- **Universal Approximation:** Cybenko (1989). “Approximation by superpositions of a sigmoidal function”

---

PART III

# Architecture Tier

---





## 🔥 Chapter 17

# Module 09: Convolutions

A convolution is a tiled matmul with structured data reuse. That reuse is what lets vision kernels hit near-peak FLOPS on hardware that would stall at fully-connected equivalents. The inductive bias of shared weights is the pedagogical story. The memory access pattern is the systems story.

### Module Info

**ARCHITECTURE TIER** | Difficulty: ●●●○ | Time: 6-8 hours | Prerequisites: 01-08

**Prerequisites — Modules 01-08.** You should have:

- Built the complete training pipeline (Modules 01-08)
- Implemented DataLoader for batch processing (Module 05)
- Comfort with parameter initialization, forward/backward passes, and optimization

If you can train an MLP on TinyDigits using your training loop and DataLoader, you're ready.

## 17.1 Overview

You have just finished the foundations tier. You can train an MLP, run optimization, move batches through a working pipeline. Welcome to the **Architecture Tier**, where the question changes from *can the network learn?* to *what should the network's structure assume about the data?*

Convolution is the answer for images. A photograph has structure that a generic MLP throws away: neighboring pixels matter together, the same edge can appear anywhere in the frame, and shifting a cat one pixel to the left should not require relearning what a cat looks like. A convolutional layer bakes those assumptions in — locality, weight sharing, and translation equivariance — and that *structural prior* is the reason CNNs dominated computer vision for a decade.

In this module you implement `Conv2d`, `MaxPool2d`, and `AvgPool2d` as explicit nested loops. No vectorization tricks, no GPU kernel calls. The loops are the point: they expose where every multiply-accumulate goes and why a single forward pass on a  $224 \times 224$  batch costs billions of operations. Once you have felt that cost in code, the optimizations real frameworks pile on top — `im2col`, `Winograd`, `cuDNN` — stop feeling like magic.

## 17.2 Learning Objectives

### By completing this module, you will:

- **Implement** `Conv2d` with explicit 7-nested loops revealing  $O(B \times C \times H \times W \times K^2 \times C_{in})$  computational complexity
- **Master** spatial dimension calculations with stride, padding, and kernel size interactions
- **Understand** receptive fields, parameter sharing, and translation equivariance in CNNs

- **Analyze** memory vs computation trade-offs: pooling reduces spatial dimensions 4x while preserving features
- **Connect** your implementations to production CNN architectures like ResNet and VGG

## 17.3 What You'll Build




Figure 17.1: **TinyTorch Spatial Operations:** Convolutional and pooling layers for hierarchical visual feature extraction.

### Implementation roadmap:

Table 17.1 lays out the implementation in order, one part at a time.

Table 17.1: **Implementation roadmap for Conv2d, MaxPool2d, and supporting layers.**

| Part | What You'll Implement            | Key Concept                            |
|------|----------------------------------|----------------------------------------|
| 1    | <code>Conv2d.__init__()</code>   | He initialization for ReLU networks    |
| 2    | <code>Conv2d.forward()</code>    | 7-nested loops for spatial convolution |
| 3    | <code>MaxPool2d.forward()</code> | Maximum selection in sliding windows   |
| 4    | <code>AvgPool2d.forward()</code> | Average pooling for smooth features    |

### The pattern you'll enable:

```
# Building a CNN block
conv = Conv2d(3, 64, kernel_size=3, padding=1)
pool = MaxPool2d(kernel_size=2, stride=2)

x = Tensor(image_batch) # (32, 3, 224, 224)
features = pool(ReLU()(conv(x))) # (32, 64, 112, 112)
```

### 17.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Dilated convolutions (PyTorch supports this with `dilation` parameter)
- Grouped convolutions (that's for efficient architectures like MobileNet)
- Depthwise separable convolutions (advanced optimization technique)

- Transposed convolutions for upsampling (used in GANs and segmentation)
- Optimized implementations (cuDNN uses Winograd algorithm and FFT convolution)

**You are building the foundational spatial operations.** Advanced convolution variants and GPU optimizations come later.

## 17.4 API Reference

This section provides a quick reference for the spatial operations you'll build. Use it as your guide while implementing and debugging convolution and pooling layers.

### 17.4.1 Conv2d Constructor

```
Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, bias=True)
```

Creates a 2D convolutional layer with learnable filters.

**Parameters:**

- `in_channels`: Number of input channels (e.g., 3 for RGB)
- `out_channels`: Number of output feature maps
- `kernel_size`: Size of convolution kernel (int or tuple)
- `stride`: Stride of convolution (default: 1)
- `padding`: Zero-padding added to input (default: 0)
- `bias`: Whether to add learnable bias (default: True)

**Weight shape:** (out\_channels, in\_channels, kernel\_h, kernel\_w)

### 17.4.2 MaxPool2d Constructor

```
MaxPool2d(kernel_size, stride=None, padding=0)
```

Creates a max pooling layer for spatial dimension reduction.

**Parameters:**

- `kernel_size`: Size of pooling window (int or tuple)
- `stride`: Stride of pooling (default: same as kernel\_size)
- `padding`: Zero-padding added to input (default: 0)

### 17.4.3 AvgPool2d Constructor

```
AvgPool2d(kernel_size, stride=None, padding=0)
```

Creates an average pooling layer for smooth spatial reduction.

**Parameters:**

- `kernel_size`: Size of pooling window (int or tuple)
- `stride`: Stride of pooling (default: same as kernel\_size)
- `padding`: Zero-padding added to input (default: 0)

### 17.4.4 BatchNorm2d Constructor

```
BatchNorm2d(num_features, eps=1e-5, momentum=0.1)
```

Batch normalization for 2D inputs (4D tensor: batch, channels, height, width). Normalizes each channel across the batch to stabilize training.

**Parameters:**

- `num_features`: Number of channels (must match input channel dimension)
- `eps`: Small constant for numerical stability (default: 1e-5)
- `momentum`: Running mean/variance update rate (default: 0.1)

### 17.4.5 Core Methods

Table 17.2 lists the core methods on the spatial layers.

Table 17.2: Core methods on spatial layers (Conv2d, MaxPool2d, BatchNorm2d).

| Method                  | Signature                                     | Description                                       |
|-------------------------|-----------------------------------------------|---------------------------------------------------|
| <code>forward</code>    | <code>forward(x: Tensor) -&gt; Tensor</code>  | Apply spatial operation to input                  |
| <code>parameters</code> | <code>parameters() -&gt; list</code>          | Return trainable parameters (Conv2d, BatchNorm2d) |
| <code>__call__</code>   | <code>__call__(x: Tensor) -&gt; Tensor</code> | Enable <code>layer(x)</code> syntax               |

### 17.4.6 Output Shape Calculation

For both convolution and pooling:

```
output_height = (input_height + 2*padding - kernel_height) ÷ stride + 1
output_width = (input_width + 2*padding - kernel_width) ÷ stride + 1
```

## 17.5 Core Concepts

This section covers the fundamental ideas you need to understand spatial operations deeply. These concepts apply to every computer vision system, from simple image classifiers to advanced object detectors.

### 17.5.1 Convolution Operation

Convolution detects local patterns by sliding a small filter (kernel) across the entire input, computing weighted sums at each position. Think of it as using a template to find matching patterns everywhere in an image.

Here's how your implementation performs this operation:

The code in ?@lst-09-convolutions-conv2d-forward makes this concrete.

```
def forward(self, x):
    # Calculate output dimensions
    out_height = (in_height + 2 * self.padding - kernel_h) // self.stride + 1
    out_width = (in_width + 2 * self.padding - kernel_w) // self.stride + 1

    # Initialize output
    output = np.zeros((batch_size, out_channels, out_height, out_width))
```

```

# Explicit 7-nested loop convolution
for b in range(batch_size):
    for out_ch in range(out_channels):
        for out_h in range(out_height):
            for out_w in range(out_width):
                in_h_start = out_h * self.stride
                in_w_start = out_w * self.stride

                conv_sum = 0.0
                for k_h in range(kernel_h):
                    for k_w in range(kernel_w):
                        for in_ch in range(in_channels):
                            input_val = padded_input[b, in_ch,
                                                    in_h_start + k_h,
                                                    in_w_start + k_w]
                            weight_val = self.weight.data[out_ch, in_ch, k_h, k_w]
                            conv_sum += input_val * weight_val

                output[b, out_ch, out_h, out_w] = conv_sum

```

: Listing 9.1 — `Conv2d.forward()` written as seven explicit nested loops, exposing the  $O(B * C_{out} * H * W * C_{in} * K^2)$  cost that `im2col` and GEMM later optimise away. {#lst-09-convolutions-conv2d-forward}

Each output pixel summarizes information from a local neighborhood in the input. A  $3 \times 3$  convolution looks at 9 pixels to produce each output value, enabling the network to detect local patterns like edges, corners, and textures.

### 17.5.2 Stride and Padding

Stride controls how far the kernel moves between positions, and padding adds zeros around the input border. Together, they determine the output spatial dimensions and receptive field coverage.

**Stride = 1** means the kernel moves one pixel at a time, producing an output nearly as large as the input. **Stride = 2** means the kernel jumps two pixels, halving the spatial dimensions and dramatically reducing computation. A stride-2 convolution processes  $4 \times$  fewer positions than stride-1.

**Padding** solves the border problem. Without padding, a  $3 \times 3$  convolution on a  $224 \times 224$  image produces a  $222 \times 222$  output, shrinking the representation. With `padding=1`, you add a 1-pixel border of zeros, keeping the output at  $224 \times 224$ . This preserves spatial dimensions and ensures edge pixels get processed as many times as center pixels.

|                       |                          |
|-----------------------|--------------------------|
| No Padding (shrinks): | Padding=1 (preserves):   |
| Input: 5x5            | Input: 5x5 → Padded: 7x7 |
| Kernel: 3x3           | Kernel: 3x3              |
| Output: 3x3           | Output: 5x5              |

|           |            |               |
|-----------|------------|---------------|
| 1 2 3 4 5 |            | 0 0 0 0 0 0 0 |
| 6 7 8 9 0 | 3x3 kernel | 0 1 2 3 4 5 0 |
| 1 2 3 4 5 | →          | 0 6 7 8 9 0 0 |
| 6 7 8 9 0 | 3x3 output | 0 1 2 3 4 5 0 |
| 1 2 3 4 5 |            | 0 6 7 8 9 0 0 |

```

0 1 2 3 4 5 0
  0 0 0 0 0 0 0

```

5x5 output preserved

The formula connecting these parameters is:

```
output_size = (input_size + 2×padding - kernel_size) / stride + 1
```

For a  $224 \times 224$  input with  $\text{kernel}=3$ ,  $\text{padding}=1$ ,  $\text{stride}=1$ :

```
output_size = (224 + 2×1 - 3) / 1 + 1 = 224
```

For the same input with  $\text{stride}=2$ :

```
output_size = (224 + 2×1 - 3) / 2 + 1 = 112
```

### 17.5.3 Receptive Fields

The **receptive field** of an output neuron is the region of the original input that can influence its value. A single  $3 \times 3$  convolution gives every output pixel a  $3 \times 3$  receptive field. Stack two  $3 \times 3$  convolutions and each output now depends on a  $5 \times 5$  region of the original input. Five stacked  $3 \times 3$  convolutions reach  $11 \times 11$ .

That growth is why depth matters. Early layers see only enough pixels to detect edges and textures. Middle layers see enough to assemble parts — eyes, wheels, corners. Deep layers see enough of the input to recognize whole objects. The hierarchy in the network mirrors the hierarchy in the data: pixels  $\rightarrow$  textures  $\rightarrow$  parts  $\rightarrow$  objects.

Receptive Field Growth:

|                               |                               |                               |
|-------------------------------|-------------------------------|-------------------------------|
| Layer 1 ( $3 \times 3$ conv): | Layer 2 ( $3 \times 3$ conv): | Layer 3 ( $3 \times 3$ conv): |
| conv):                        |                               |                               |
| $\rightarrow 3 \times 3$ RF   | $\rightarrow 5 \times 5$ RF   | $\rightarrow 7 \times 7$ RF   |

Stacking  $N$   $3 \times 3$  convolutions:

Receptive Field =  $1 + 2 \times N$

VGG-16 uses this principle: stack many small kernels instead of few large ones.

Parameter sharing means the same  $3 \times 3$  kernel processes every position in the image. This drastically reduces parameters compared to fully connected layers while maintaining translation equivariance: if you shift the input, the output shifts identically.

### 17.5.4 Pooling Operations

Pooling reduces spatial dimensions while preserving important features. Max pooling selects the strongest activation in each window, preserving sharp features like edges. Average pooling computes the mean, creating smoother, more general features.

Here's how max pooling works in your implementation:

The code in `?@lst-09-convolutions-maxpool2d-forward` makes this concrete.

```

def forward(self, x):
    # Calculate output dimensions
    out_height = (in_height + 2 * self.padding - kernel_h) // self.stride + 1
    out_width = (in_width + 2 * self.padding - kernel_w) // self.stride + 1

    output = np.zeros((batch_size, channels, out_height, out_width))

```

```

# Explicit nested loop max pooling
for b in range(batch_size):
    for c in range(channels):
        for out_h in range(out_height):
            for out_w in range(out_width):
                in_h_start = out_h * self.stride
                in_w_start = out_w * self.stride

                # Find maximum in window
                max_val = -np.inf
                for k_h in range(kernel_h):
                    for k_w in range(kernel_w):
                        input_val = padded_input[b, c,
                                                in_h_start + k_h,
                                                in_w_start + k_w]
                        max_val = max(max_val, input_val)

                output[b, c, out_h, out_w] = max_val

```

: Listing 9.2 — `MaxPool2d.forward()` selects the largest value in each sliding window, no learnable parameters — just comparisons. {#lst-09-convolutions-maxpool2d-forward}

A  $2 \times 2$  max pooling with `stride=2` divides spatial dimensions by 2, reducing memory and computation by  $4\times$ . For a  $224 \times 224 \times 64$  feature map (12.3 MB in float32), pooling produces  $112 \times 112 \times 64$  (3.1 MB), saving 9.2 MB per layer.

Max pooling provides translation invariance: if a cat's ear moves one pixel, the max in that region remains roughly the same, making the network robust to small shifts. This is crucial for object recognition where precise pixel alignment doesn't matter.

Average pooling smooths features by averaging windows, useful for global feature summarization. Modern architectures often use global average pooling (averaging entire feature maps to single values) instead of fully connected layers, dramatically reducing parameters.

### 17.5.5 Output Shape Calculation

Understanding output shapes is critical for building CNNs. A shape mismatch crashes your network, while correct dimensions ensure features flow properly through layers.

The output shape formula applies to both convolution and pooling:

$$H_{out} = (H_{in} + 2 \times padding - kernel\_h) / stride + 1$$

$$W_{out} = (W_{in} + 2 \times padding - kernel\_w) / stride + 1$$

The floor operation (`||`) ensures integer dimensions. If the calculation doesn't divide evenly, the rightmost and bottommost regions get ignored.

**Example calculations:**

Input: (32, 3, 224, 224) [batch=32, RGB channels,  $224 \times 224$  image]

Conv2d(3, 64, kernel\_size=3, padding=1, stride=1):  $H_{out} = (224 + 2 \times 1 - 3) / 1 + 1 = 224$   $W_{out} = (224 + 2 \times 1 - 3) / 1 + 1 = 224$  Output: (32, 64, 224, 224)

MaxPool2d(kernel\_size=2, stride=2):  $H_{out} = (224 + 0 - 2) / 2 + 1 = 112$   $W_{out} = (224 + 0 - 2) / 2 + 1 = 112$  Output: (32, 64, 112, 112)

Conv2d(64, 128, kernel\_size=3, padding=0, stride=2):  $H_{out} = (112 + 0 - 3) / 2 + 1 = 55$   $W_{out} = (112 + 0 - 3) / 2 + 1 = 55$  Output: (32, 128, 55, 55)

**Common patterns:** - **Same convolution** (padding=1, stride=1, kernel=3): Preserves spatial dimensions  
 - **Stride-2 convolution:** Halves dimensions, replaces pooling in some architectures (ResNet) - **2×2 pooling, stride=2:** Classic dimension reduction, halves H and W

### 17.5.6 Computational Complexity

Convolution is expensive. The explicit loops reveal exactly why: you're visiting every position in the output, and for each position, sliding over the entire kernel across all input channels.

For a single Conv2d forward pass:

Operations =  $B \times C_{out} \times H_{out} \times W_{out} \times C_{in} \times K_h \times K_w$

**Example:** Batch=32, Input=(3, 224, 224), Conv2d(3→64, kernel=3, padding=1, stride=1)

Operations =  $32 \times 64 \times 224 \times 224 \times 3 \times 3 \times 3 = 32 \times 64 \times \{\text{python}\} \text{ complexity\_hw} \times \{\text{python}\} \text{ complexity\_kernel} = \{\text{python}\} \text{ complexity\_ops}$  multiply-accumulate operations  $\approx \{\text{python}\} \text{ complexity\_approx billion operations}$  per forward pass!

Executing billions of scattered operations sequentially via deeply nested `for` loops yields catastrophic cache-miss rates because it fails to leverage **temporal locality** (reusing recently accessed data) and **spatial locality** (accessing contiguous memory blocks). This forces engineers to radically restructure how convolutions map to massively parallel hardware.

#### Systems Implication: Memory Duplication (im2col)

Modern GPUs physically cannot execute branching, nested `for` loops efficiently; however, they possess thousands of highly specialized cores explicitly optimized for dense General Matrix Multiplication (GEMM) via hardware cache tiling. To bridge this architectural gap, frameworks employ `im2col` (Image to Column). This algorithm systematically duplicates the spatial image memory in VRAM, unrolling the overlapping sliding windows into one massive, contiguous 2D matrix. By deliberately sacrificing raw memory capacity, `im2col` artificially creates perfect **spatial locality** for the GEMM hardware. This structural reorganization ensures that cache tiling can maximally exploit **temporal locality**, transforming a disjointed spatial traversal into a monolithic matrix multiplication and achieving a 100× throughput acceleration.

While `im2col` elegantly bypasses the compute bottleneck, it aggressively transfers that pressure directly onto VRAM capacity. This harsh tension between tensor memory duplication and computational vectorization underscores why spatial kernel size is aggressively optimized. Because a  $7 \times 7$  kernel imposes `{python} kernel_ratio`× more compute and spatial redundancy than a  $3 \times 3$  kernel, modern architectures (like VGG and ResNet) strictly favor stacking deep sequences of  $3 \times 3$  convolutions rather than deploying monolithic, large-footprint kernels.

Pooling operations are cheap by comparison: no learnable parameters, just comparison or addition operations. A  $2 \times 2$  max pooling visits each output position once and compares 4 values, requiring only  $4 \times$  comparisons per output.

Table 17.3 summarises the time and memory cost of the core operations.

Table 17.3: **Computational complexity of 2D spatial operations.**

| Operation                  | Complexity                                                       | Notes                                      |
|----------------------------|------------------------------------------------------------------|--------------------------------------------|
| Conv2d ( $K \times K$ )    | $O(B \times C_{out} \times H \times W \times C_{in} \times K^2)$ | Cubic in spatial dims, quadratic in kernel |
| MaxPool2d ( $K \times K$ ) | $O(B \times C \times H \times W \times K^2)$                     | No channel mixing, just spatial reduction  |



| Operation       | Complexity                                   | Notes                             |
|-----------------|----------------------------------------------|-----------------------------------|
| AvgPool2d (K×K) | $O(B \times C \times H \times W \times K^2)$ | Same as MaxPool but with addition |

Memory consumption follows the output shape. A (32, 64, 224, 224) float32 tensor requires:

$32 \times 64 \times 224 \times 224 \times 4 \text{ bytes} = 392 \text{ MB}$

That single tensor is your batch's activations after one layer. Doubling batch size doubles that memory. GPUs have limited memory (typically 8-24 GB), which is what ultimately caps batch size and feature-map resolution in practice.

## 17.6 Common Errors

These are the errors you'll encounter most often when working with spatial operations. Understanding why they happen will save you hours of debugging CNNs.

### 17.6.1 Shape Mismatch in Conv2d

**Error:** ValueError: Expected 4D input (batch, channels, height, width), got (3, 224, 224)

Conv2d requires 4D input: (batch, channels, height, width). If you forget the batch dimension, the layer interprets channels as batch, height as channels, causing chaos.

**Fix:** Add batch dimension: `x = x.reshape(1, 3, 224, 224)` or ensure your data pipeline always includes batch dimension.

### 17.6.2 Dimension Calculation Errors

**Error:** Output shape is 55 when you expected 56

The floor operation in output dimension calculation can surprise you. If  $(\text{input} + 2 \times \text{padding} - \text{kernel}) / \text{stride}$  doesn't divide evenly, the result gets floored.

**Example:**

Input:  $224 \times 224$ , kernel=3, padding=0, stride=2 output\_size =  $(224 + 0 - 3) // 2 + 1 = 221 // 2 + 1 = 110 + 1 = 111$

**Fix:** Use calculators or test with dummy data to verify dimensions before building full architecture.

### 17.6.3 Padding Value Confusion

**Error:** Max pooling produces zeros at borders when using `padding > 0`

If you pad max pooling input with zeros (`constant_values=0`), and your feature map has negative values, the padded zeros will be selected as maximums at borders, creating artifacts.

**Fix:** Pad max pooling with `-np.inf`:

```
padded_input = np.pad(x.data, ..., constant_values=-np.inf)
```

### 17.6.4 Stride/Kernel Mismatch in Pooling

**Error:** Overlapping pooling windows when `stride ≠ kernel_size`

By convention, pooling uses non-overlapping windows: `stride = kernel_size`. If you accidentally set `stride=1` with `kernel=2`, windows overlap, creating redundant computation and unexpected behavior.

**Fix:** Ensure `stride = kernel_size` for pooling, or set `stride=None` to use default (equals `kernel_size`).

### 17.6.5 Memory Overflow

**Error:** `RuntimeError: CUDA out of memory` or system hangs

Large feature maps consume enormous memory. A batch of 64 images at  $224 \times 224 \times 64$  channels uses 0.8 GB for a single layer's output. Stack fifty layers and the activations alone — needed for backprop — outgrow most GPUs.

**Fix:** Reduce batch size, use smaller images, or add more pooling layers to reduce spatial dimensions faster.

## 17.7 Production Context

### 17.7.1 Your Implementation vs. PyTorch

Your TinyTorch spatial operations and PyTorch's `torch.nn.Conv2d` share the same conceptual foundation: sliding kernels, stride, padding, output shape formulas. The differences lie in optimization and hardware support.

Table 17.4 places your implementation side by side with the production reference for direct comparison.

Table 17.4: Feature comparison between TinyTorch Conv2d and PyTorch cuDNN.

| Feature             | Your Implementation        | PyTorch                                       |
|---------------------|----------------------------|-----------------------------------------------|
| <b>Backend</b>      | NumPy loops (Python)       | cuDNN (CUDA C++)                              |
| <b>Speed</b>        | 1x (baseline)              | 100-1000x faster on GPU                       |
| <b>Optimization</b> | Explicit loops             | im2col + GEMM, Winograd, FFT                  |
| <b>Memory</b>       | Straightforward allocation | Memory pooling, gradient checkpointing        |
| <b>Features</b>     | Basic conv + pool          | Dilated, grouped, transposed, 3D convolutions |

### 17.7.2 Code Comparison

The following comparison shows equivalent operations in TinyTorch and PyTorch. Notice how the API mirrors perfectly, making your knowledge transfer directly to production frameworks.

## 17.8 Your TinyTorch

```
from tinytorch.core.spatial import Conv2d, MaxPool2d, AvgPool2d
from tinytorch.core.activations import ReLU

# Build a CNN block
conv1 = Conv2d(3, 64, kernel_size=3, padding=1)
conv2 = Conv2d(64, 128, kernel_size=3, padding=1)
pool = MaxPool2d(kernel_size=2, stride=2) # Or use AvgPool2d for smooth features
relu = ReLU()
```

```
# Forward pass
x = Tensor(image_batch) # (32, 3, 224, 224)
x = relu(conv1(x))      # (32, 64, 224, 224)
x = pool(x)             # (32, 64, 112, 112)
x = relu(conv2(x))      # (32, 128, 112, 112)
x = pool(x)             # (32, 128, 56, 56)
```

## 17.9 PyTorch

```
import torch
import torch.nn as nn

# Build a CNN block
conv1 = nn.Conv2d(3, 64, kernel_size=3, padding=1)
conv2 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
pool = nn.MaxPool2d(kernel_size=2, stride=2)
relu = nn.ReLU()

# Forward pass (identical structure!)
x = torch.tensor(image_batch, dtype=torch.float32)
x = relu(conv1(x))
x = pool(x)
x = relu(conv2(x))
x = pool(x)
```

Let's walk through each line to understand the comparison:

- **Lines 1-2 (Imports):** TinyTorch separates spatial operations and activations into different modules; PyTorch consolidates into `torch.nn`. Same concepts, different organization.
- **Lines 4-7 (Layer creation):** Identical API. Both use `Conv2d(in, out, kernel_size, padding)` and `MaxPool2d(kernel_size, stride)`. The parameter names and semantics are identical.
- **Line 10 (Input):** TinyTorch wraps in `Tensor`; PyTorch uses `torch.tensor()` with explicit dtype. Same abstraction.
- **Lines 11-14 (Forward pass):** Identical call patterns. ReLU activations, pooling for dimension reduction, growing channels (3→64→128). This is the standard CNN building block.
- **Shapes:** Every intermediate shape matches between frameworks because the formulas are identical.

### 💡 What's Identical

Convolution mathematics, stride and padding formulas, receptive field growth, and parameter sharing. The APIs are intentionally identical so your understanding transfers directly to production systems.

### 17.9.1 Why Spatial Operations Matter at Scale

The scale of production vision systems is what makes convolution optimization a first-class engineering concern:

- **ResNet-50** — 25 million parameters, **4 billion operations** per image, thousands of images per second in serving
- **YOLO object detection** — 30 FPS video at 1080p, **~60 billion convolution operations per second**
- **Self-driving cars** — 10+ CNN models running simultaneously across 6 cameras at 30 FPS, **~300 billion operations per second** inside a 50ms latency budget

Your educational Conv2d takes hundreds of milliseconds on CPU for a single forward pass. The equivalent PyTorch call runs in single-digit milliseconds on GPU, and the gap is bridged by three ideas you should know by name:

- **im2col + GEMM** — reshape the convolution into a matrix multiply and hand it to a tuned BLAS kernel
- **Winograd** — algebraic identity that cuts the multiplication count for  $3 \times 3$  kernels by  $\sim 2.25 \times$
- **FFT convolution** — for large kernels, Fourier-domain multiplication trades  $O(n^2)$  for  $O(n \log n)$

Module 17 is where you replace your loops with the first of these. Before you can optimize convolution, you have to feel why it is slow — which is exactly what the loops you wrote in this module make undeniable.

## 17.10 Check Your Understanding

### 💡 Check Your Understanding — Convolutions

Before moving on, verify you can articulate each of the following:

- The im2col trick — how it transforms a convolution into a single GEMM at the cost of memory duplication, and why that cost is worth paying on hardware that is optimized for dense matrix multiply.
- Why weight sharing gives convolution a  $\sim 5,000 \times$  parameter advantage over a dense layer on the same image, and why that is a structural — not incidental — property.
- How to compute output shape, parameter count, and MACs for any `Conv2d(C_in, C_out, K, stride, padding)` configuration without reaching for a calculator.
- Why receptive field grows with depth (and with stride products), and why stacking  $3 \times 3$  convolutions beats a single  $7 \times 7$  — both in parameters and in the FLOPs per covered region.

If any of these feels fuzzy, revisit the Computational Complexity and Receptive Fields sections before moving on.

Use these questions to test the systems intuition — output shapes, parameters, and computational cost — that you'll need every time you reach for a real CNN architecture.

#### Q1: Output Shape Calculation

Given input (32, 3, 128, 128), what's the output shape after `Conv2d(3, 64, kernel_size=5, padding=2, stride=2)`?

### 💡 Answer

Calculate height and width:

$$H_{\text{out}} = (128 + 2 \times 2 - 5) / 2 + 1 = (128 + 4 - 5) / 2 + 1 = 127 / 2 + 1 = 63 + 1 = 64$$

$$W_{\text{out}} = (128 + 2 \times 2 - 5) / 2 + 1 = 64$$

Output shape: (32, 64, 64, 64)

Batch stays the same, channels jump 3→64, and spatial dimensions halve because stride=2.

#### Q2: Parameter Counting

How many parameters in `Conv2d(3, 64, kernel_size=3, bias=True)`?

 **Answer**

Weight parameters:  $\text{out\_channels} \times \text{in\_channels} \times \text{kernel\_h} \times \text{kernel\_w}$

Weight:  $64 \times 3 \times 3 \times 3 = 1,728$  parameters Bias: 64 parameters Total: **1,792 parameters**

Compare this to a fully connected layer for  $224 \times 224$  RGB images:

$\text{Dense}(224 \times 224 \times 3, 64) = 150,528 \times 64 = 9,633,792$  parameters.

Convolution uses **5,376× fewer parameters** for the same task — that is the price of weight sharing, and it is the structural advantage that makes deep CNNs trainable on commodity hardware.

**Q3: Computational Complexity**

For input (16, 64, 56, 56) and `Conv2d(64, 128, kernel_size=3, padding=1, stride=1)`, how many multiply-accumulate operations?

 **Answer**

Operations =  $B \times C_{\text{out}} \times H_{\text{out}} \times W_{\text{out}} \times C_{\text{in}} \times K_{\text{h}} \times K_{\text{w}}$

First calculate output dimensions:

$H_{\text{out}} = (56 + 2 \times 1 - 3) / 1 + 1 = 56$   $W_{\text{out}} = (56 + 2 \times 1 - 3) / 1 + 1 = 56$

Then total operations:

$16 \times 128 \times 56 \times 56 \times 64 \times 3 \times 3 = 16 \times 128 \times 3,136 \times 576 = 3,699,376,128$  operations  $\approx 3.7$  billion operations per forward pass.

This is why batch size directly impacts training time: doubling the batch doubles the operations.

**Q4: Memory Calculation**

What's the memory requirement for storing the output of `Conv2d(3, 256, kernel_size=7, stride=2, padding=3)` on input (64, 3, 224, 224)?

 **Answer**

First calculate output dimensions:

$H_{\text{out}} = (224 + 2 \times 3 - 7) / 2 + 1 = (224 + 6 - 7) / 2 + 1 = 223 / 2 + 1 = 111 + 1 = 112$   $W_{\text{out}} = 112$

Output shape: (64, 256, 112, 112)

Memory (float32 = 4 bytes):

$64 \times 256 \times 112 \times 112 \times 4 = 822,083,584$  bytes  $\approx$  **784 MB** for a single layer's output.

Deep CNNs need GPUs with large memory (16 GB and up) for exactly this reason. Backprop must keep every layer's activations resident, and fifty layers of this size will not fit anywhere else.

**Q5: Receptive Field Growth**

Starting with  $224 \times 224$  input, you stack: `Conv(3x3, stride=1)` → `MaxPool(2x2, stride=2)` → `Conv(3x3, stride=1)` → `Conv(3x3, stride=1)`. What's the receptive field of the final layer?

 **Answer**

Track receptive field growth through each layer:

Layer 1 — `Conv(3x3, stride=1)`: RF = 3 Layer 2 — `MaxPool(2x2, stride=2)`: RF =  $3 + (2-1) \times 1 = 4$

Layer 3 — `Conv(3x3, stride=1)`: RF =  $4 + (3-1) \times 2 = 8$  (stride accumulates) Layer 4 — `Conv(3x3, stride=1)`: RF =  $8 + (3-1) \times 2 = 12$

**Receptive field = 12x12**

Each neuron in the final layer sees a  $12 \times 12$  region of the original input. This is why stacking layers with stride or pooling matters: it grows the receptive field so deeper layers can detect larger patterns.

Formula:  $RF_{new} = RF_{old} + (kernel\_size - 1) \times stride\_product$   
 where  $stride\_product$  is the accumulated stride from all previous layers.

## 17.11 Key Takeaways

- **Convolution is a structural prior, not just an operation:** Locality, weight sharing, and translation equivariance are baked in — those priors are the reason CNNs beat MLPs on images with orders of magnitude fewer parameters.
- **The 7-nested loop exposes the cost:**  $O(B \times C_{out} \times H \times W \times C_{in} \times K^2)$  is cubic in spatial dimensions and quadratic in kernel size, which is why a single ResNet-50 forward pass on one image is ~4 billion MACs.
- **im2col trades VRAM for throughput:** Unrolling sliding windows into one dense matrix hands the work to GPU GEMM kernels that exploit cache tiling. The duplicated memory is the tax you pay for  $\sim 100\times$  acceleration.
- **Output shape, parameter count, and receptive field are the three CNN formulas you live by:**  $(H + 2p - K)/s + 1$ ,  $C_{out} \times C_{in} \times K^2$ , and  $RF_{new} = RF_{old} + (K-1) \times stride\_product$ . Every architectural decision traces back to these.
- **Activation memory, not parameter memory, caps batch size:** A single (64, 256, 112, 112) feature map is ~784 MB. Fifty such layers, times two for backward, is what forces gradient checkpointing in practice.

**Coming next:** Convolution is the structural prior for grids of pixels. Text has no grid — it is a discrete sequence with no fixed alphabet. Module 10 builds BPE and WordPiece tokenizers and introduces the first real design question of the language stack: vocabulary.

## 17.12 Further Reading

For students who want to understand the academic foundations and explore spatial operations further:

### 17.12.1 Seminal Papers

- **Gradient-Based Learning Applied to Document Recognition** - LeCun et al. (1998). The paper that launched convolutional neural networks, introducing LeNet-5 for handwritten digit recognition. Essential reading for understanding why convolution works for vision.
  - **Systems Implication:** Weight sharing in convolutions drastically reduced the required parameter count compared to dense, fully connected layers. This architectural breakthrough made high-fidelity inference mathematically feasible on the severely constrained memory and cache architecture of 1990s-era CPUs. [IEEE](#)
- **ImageNet Classification with Deep Convolutional Neural Networks** - Krizhevsky et al. (2012). AlexNet, the breakthrough that demonstrated CNNs could conquer the massive ImageNet dataset, cementing ReLU, dropout, and data augmentation into the modern deep learning canon.
  - **Systems Implication:** The parameter density of AlexNet exceeded the physical limits of a single GTX 580 GPU (3GB VRAM). The authors were forced to shard the layer weights and feature maps across two independent GPUs, unintentionally giving birth to the modern paradigm of Model Parallelism. [NeurIPS](#)
- **Very Deep Convolutional Networks for Large-Scale Image Recognition** - Simonyan & Zisserman (2014). VGG networks proved that relentlessly stacking many compact  $3\times 3$  convolutions yields vastly superior hierarchical features compared to using sparse, large kernels.

- **Systems Implication:** While VGG's uniform architecture simplified GPU kernel scheduling, its decision to maintain high-resolution, multi-channel feature maps deep into the network created a catastrophic memory footprint, severely constraining batch sizes and dictating the need for more efficient downstream architectures. [arXiv:1409.1556](#)
- **Deep Residual Learning for Image Recognition** - He et al. (2015). ResNet introduced residual skip connections, enabling the stable training of ultra-deep, 100+ layer networks and entirely conquering the vanishing gradient problem.
  - **Systems Implication:** Residual skip connections structurally necessitated keeping early-layer spatial activations alive in VRAM long after their immediate computation finished. This elongated tensor lifecycle fundamentally complicated GPU memory allocator design and severely worsened Activation Pinning during the forward pass. [arXiv:1512.03385](#)

### 17.12.2 Additional Resources

- **CS231n: Convolutional Neural Networks for Visual Recognition** - Stanford course notes with excellent visualizations of convolution, receptive fields, and feature maps: <https://cs231n.github.io/convolutional-networks/>
- **Textbook:** “Deep Learning” by Goodfellow, Bengio, and Courville - Chapter 9 covers convolutional networks with mathematical depth
- **Distill.pub:** “Feature Visualization” - Interactive article showing what CNN filters learn at different depths: <https://distill.pub/2017/feature-visualization/>

## 17.13 What's Next

### **i** Coming Up — Module 10: Tokenization

Convolution is the structural prior for grids of pixels. The next data type — text — is not a grid. It is a discrete sequence with no fixed alphabet, no fixed length, and no notion of intensity. Before any model can read English you have to answer a more basic question: how do you turn a string into numbers a network can multiply against? In Module 10 you implement BPE and WordPiece tokenizers and meet the first real architectural decision of the language stack: vocabulary.

#### How your spatial operations carry forward:

Table 17.5 traces how this module is reused by later parts of the curriculum.

Table 17.5: How spatial ops feed into later milestone and optimization modules.

| Module                         | What it does                   | Your spatial ops in action                               |
|--------------------------------|--------------------------------|----------------------------------------------------------|
| <b>Milestone 3: CNN</b>        | Complete CNN for CIFAR-10      | Stack your Conv2d and MaxPool2d for image classification |
| <b>Module 17: Acceleration</b> | Optimize convolution           | Replace loops with im2col and vectorized GEMM            |
| <b>Vision Projects</b>         | Object detection, segmentation | Your spatial foundations scale to advanced architectures |

## 17.14 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## 🔥 Chapter 18

# Module 10: Tokenization

A language model never sees text. It sees token IDs. The vocabulary you pick upstream silently caps sequence length, embedding-table size, and every attention cost downstream. That single design choice is the first systems decision of the language stack.

### Module Info

**ARCHITECTURE TIER** | Difficulty: ●●○○ | Time: 4-6 hours | Prerequisites: 01-08

You should have completed the Foundation tier:

- Tensor operations (Module 01)
- Basic neural network components (Modules 02-04)
- Training fundamentals (Modules 05-07)

Tokenization stands largely on its own — it works with strings and dictionaries, not gradients. If you can manipulate Python strings, you're ready.

## 18.1 Overview

Text isn't the input to a language model. *Tokens* are. Before a single matrix multiply happens inside GPT, every character of your prompt has already been chopped, merged, and looked up in a fixed vocabulary — and that one upstream choice silently decides how long your sequences are, how big your embedding table is, and how much an inference call costs.

In this module you build two tokenizers from scratch: a character-level tokenizer (one character, one token) and a Byte Pair Encoding (BPE) tokenizer that learns subword units from data. Doing both surfaces the central trade-off: small vocabularies yield long sequences and tiny embedding tables; large vocabularies yield short sequences but multi-megabyte embeddings. Attention cost scales quadratically with sequence length, so this trade is rarely close.

By the end you can explain why GPT uses ~50,000 tokens, how tokenizers degrade gracefully on unknown words, and how the vocabulary you pick today caps the throughput of every model you train tomorrow.

## 18.2 Learning Objectives

### By completing this module, you will:

- **Implement** character-level tokenization for robust text coverage and BPE tokenization for efficient subword representation
- **Understand** the vocabulary size versus sequence length trade-off and its impact on memory and computation
- **Master** encoding and decoding operations that convert between text and numerical token IDs

- **Connect** your implementation to production tokenizers used in GPT, BERT, and modern language models

## 18.3 What You'll Build




Figure 18.1: **TinyTorch Tokenization Infrastructure:** Converting raw text into model-ready numerical sequences.

### Implementation roadmap:

Table 18.1 lays out the implementation in order, one part at a time.

Table 18.1: **Implementation roadmap for the Tokenizer classes.**

| Part | What You'll Implement | Key Concept                                     |
|------|-----------------------|-------------------------------------------------|
| 1    | Tokenizer base class  | Interface contract: encode/decode               |
| 2    | CharTokenizer         | Character-level vocabulary, perfect coverage    |
| 3    | BPETokenizer          | Byte Pair Encoding, learning merges             |
| 4    | Vocabulary building   | Unique character extraction, frequency analysis |
| 5    | Utility functions     | Dataset processing, analysis tools              |

### The pattern you'll enable:

```
# Converting text to numbers for neural networks
tokenizer = BPETokenizer(vocab_size=1000)
tokenizer.train(corpus)
```

```
token_ids = tokenizer.encode("Hello world")  # [142, 1847, 2341]
```

### 18.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- GPU-accelerated tokenization (production tokenizers use Rust/C++)
- Advanced segmentation algorithms (SentencePiece, Unigram models)
- Language-specific preprocessing (Unicode normalization, byte-level fallback)
- Tokenizer serialization and loading (PyTorch handles this with `save_pretrained()`)

**You are building the conceptual foundation.** Production optimizations come later.

## 18.4 API Reference

This section provides a quick reference for the tokenization classes you'll build. Think of it as your cheat sheet while implementing and debugging.

### 18.4.1 Base Tokenizer Interface

```
Tokenizer()
```

- Abstract base class defining the tokenizer contract
- All tokenizers must implement `encode()` and `decode()`

### 18.4.2 CharTokenizer

```
CharTokenizer(vocab: Optional[List[str]] = None)
```

- Character-level tokenizer treating each character as a token
- `vocab`: Optional list of characters to include in vocabulary

Table 18.2 lists the methods you will implement on this class.

Table 18.2: Methods on the CharTokenizer class.

| Method                   | Signature                                              | Description                           |
|--------------------------|--------------------------------------------------------|---------------------------------------|
| <code>build_vocab</code> | <code>build_vocab(corpus: List[str]) -&gt; None</code> | Extract unique characters from corpus |
| <code>encode</code>      | <code>encode(text: str) -&gt; List[int]</code>         | Convert text to character IDs         |
| <code>decode</code>      | <code>decode(tokens: List[int]) -&gt; str</code>       | Convert character IDs back to text    |

**Properties:** - `vocab`: List of characters in vocabulary - `vocab_size`: Total number of unique characters + special tokens - `char_to_id`: Mapping from characters to IDs - `id_to_char`: Mapping from IDs to characters - `unk_id`: ID for unknown characters (always 0)

### 18.4.3 BPETokenizer

```
BPETokenizer(vocab_size: int = 1000)
```

- Byte Pair Encoding tokenizer learning subword units
- `vocab_size`: Target vocabulary size after training

Table 18.3 lists the methods you will implement on this class.

Table 18.3: Methods on the BPETokenizer class.

| Method              | Signature                                                                | Description                       |
|---------------------|--------------------------------------------------------------------------|-----------------------------------|
| <code>train</code>  | <code>train(corpus: List[str], vocab_size: int = None) -&gt; None</code> | Learn BPE merges from corpus      |
| <code>encode</code> | <code>encode(text: str) -&gt; List[int]</code>                           | Convert text to subword token IDs |
| <code>decode</code> | <code>decode(tokens: List[int]) -&gt; str</code>                         | Convert token IDs back to text    |

#### Helper Methods:

Table 18.4 lists the methods you will implement on this class.

Table 18.4: Internal helper methods used by the BPETokenizer.

| Method                        | Signature                                                                  | Description                                            |
|-------------------------------|----------------------------------------------------------------------------|--------------------------------------------------------|
| <code>_get_word_tokens</code> | <code>_get_word_tokens(word: str) -&gt; List[str]</code>                   | Convert word to character list with end-of-word marker |
| <code>_get_pairs</code>       | <code>_get_pairs(word_tokens: List[str]) -&gt; Set[Tuple[str, str]]</code> | Extract all adjacent character pairs                   |
| <code>_apply_merges</code>    | <code>_apply_merges(tokens: List[str]) -&gt; List[str]</code>              | Apply learned merge rules to token sequence            |
| <code>_build_mappings</code>  | <code>_build_mappings() -&gt; None</code>                                  | Build token-to-ID and ID-to-token dictionaries         |

**Properties:** - `vocab`: List of tokens (characters + learned merges) - `vocab_size`: Total vocabulary size - `merges`: List of learned merge rules (pair tuples) - `token_to_id`: Mapping from tokens to IDs - `id_to_token`: Mapping from IDs to tokens

### 18.4.4 Utility Functions

Table 18.5 lists the utility helpers that accompany the tokenizers.

Table 18.5: Utility functions for tokenizer creation and analysis.

| Function                          | Signature                                                                                                    | Description                           |
|-----------------------------------|--------------------------------------------------------------------------------------------------------------|---------------------------------------|
| <code>create_tokenizer</code>     | <code>create_tokenizer(strategy: str, vocab_size: int, corpus: List[str]) -&gt; Tokenizer</code>             | Factory for creating tokenizers       |
| <code>tokenize_dataset</code>     | <code>tokenize_dataset(texts: List[str], tokenizer: Tokenizer, max_length: int) -&gt; List[List[int]]</code> | Batch tokenization with length limits |
| <code>analyze_tokenization</code> | <code>analyze_tokenization(texts: List[str], tokenizer: Tokenizer) -&gt; Dict[str, float]</code>             | Compute statistics and metrics        |

## 18.5 Core Concepts

This section covers the fundamental ideas you need to understand tokenization deeply. These concepts apply to every NLP system, from simple chatbots to large language models.

### 18.5.1 Text to Numbers

Neural networks process numbers, not text. When you pass the string “Hello” to a model, it must first become a sequence of integers. This transformation happens in four steps: split text into tokens (units of meaning), build a vocabulary mapping each unique token to an integer ID, encode text by looking up each token’s ID, and enable decoding to reconstruct the original text from IDs.

The simplest approach treats each character as a token. Consider the word “hello”: split into characters `['h', 'e', 'l', 'l', 'o']`, build a vocabulary with IDs `{'h': 1, 'e': 2, 'l': 3, 'o': 4}`, encode to `[1, 2, 3, 3, 4]`, and decode back by reversing the lookup. This implementation is beautifully simple:

```
def encode(self, text: str) -> List[int]:
    """Encode text to list of character IDs."""
    tokens = []
    for char in text:
        tokens.append(self.char_to_id.get(char, self.unk_id))
    return tokens
```

The elegance is in the simplicity: iterate through each character, look up its ID in the vocabulary dictionary, and use the unknown token ID for unseen characters. This gives perfect coverage: any text can be encoded without errors, though the sequences can be long.

### 18.5.2 Vocabulary Building

Before encoding text, you need a vocabulary: the complete set of tokens your tokenizer recognizes. For character-level tokenization, this means extracting all unique characters from a training corpus.

To construct this vocabulary systematically, we extract every unique character observed in the training corpus and sort them to ensure a consistent mapping:

The code in `?lst-10-tokenization-build-vocab` makes this concrete.

```
def build_vocab(self, corpus: List[str]) -> None:
    """Build vocabulary from a corpus of text."""
    # Collect all unique characters
    all_chars = set()
    for text in corpus:
        all_chars.update(text)

    # Sort for consistent ordering
    unique_chars = sorted(list(all_chars))

    # Rebuild vocabulary with <UNK> token first
    self.vocab = ['<UNK>'] + unique_chars
    self.vocab_size = len(self.vocab)

    # Rebuild mappings
    self.char_to_id = {char: idx for idx, char in enumerate(self.vocab)}
    self.id_to_char = {idx: char for idx, char in enumerate(self.vocab)}
```

: Listing 10.1 — `CharTokenizer.build_vocab()` collects unique characters, reserves index 0 for `<UNK>`, and builds the forward and reverse ID maps. {#lst-10-tokenization-build-vocab}

The special `<UNK>` token at position 0 handles characters not in the vocabulary. When encoding text with unknown characters, they all map to ID 0. This graceful degradation prevents crashes while signaling that information was lost.

Character vocabularies are tiny: typically 50-200 tokens depending on language, which means small embedding tables. A 100-character vocabulary with 512-dimensional embeddings requires only **51,200 parameters**, about **200 KB** of memory — dramatically smaller than word-level vocabularies with 100,000+ entries.

### 18.5.3 Byte Pair Encoding (BPE)

Character tokenization has one fatal flaw: the sequences are too long. A 50-word sentence becomes ~250 tokens, attention costs scale with the *square* of that length, and the model has to learn from scratch that `'h', 'e', 'l', 'l', 'o'` is the same thing as `'h', 'e', 'l', 'p'` plus one swap.

BPE fixes this by learning subword units from the data itself. The algorithm is four lines: start with a character-level vocabulary, count all adjacent character pairs in the corpus, merge the most frequent pair into a new token, and repeat until you hit the target vocabulary size.

Consider training BPE on the corpus `["hello", "hello", "help"]`. Each word starts with end-of-word markers: `['h', 'e', 'l', 'l', 'o', '</w>']`, `['h', 'e', 'l', 'l', 'o', '</w>']`, `['h', 'e', 'l', 'p', '</w>']`. Count all pairs: `('h', 'e')` appears 3 times, `('e', 'l')` appears 3 times, `('l', 'l')` appears 2 times. The most frequent is `('h', 'e')`, so merge it:

```
# Merge operation: ('h', 'e') -> 'he'
# Before:
['h', 'e', 'l', 'l', 'o', '</w>'] -> ['he', 'l', 'l', 'o', '</w>']
['h', 'e', 'l', 'l', 'o', '</w>'] -> ['he', 'l', 'l', 'o', '</w>']
['h', 'e', 'l', 'p', '</w>'] -> ['he', 'l', 'p', '</w>']
```

The vocabulary grows from `['h', 'e', 'l', 'l', 'o', 'p', '</w>']` to `['h', 'e', 'l', 'l', 'o', 'p', '</w>', 'he']`. Continue merging: next most frequent is `('l', 'l')`, so merge to get `'ll'`. The vocabulary becomes `['h', 'e', 'l', 'l', 'o', 'p', '</w>', 'he', 'll']`. After sufficient merges, “hello” encodes as `['he', 'll', 'o', '</w>']` (3 tokens instead of 5 characters).

Here's how the training loop works:

The code in `#lst-10-tokenization-bpe-train` makes this concrete.

```
while len(self.vocab) < self.vocab_size:
    # Count all pairs across all words
    pair_counts = Counter()
    for word, freq in word_freq.items():
        tokens = word_tokens[word]
        pairs = self._get_pairs(tokens)
        for pair in pairs:
            pair_counts[pair] += freq

    if not pair_counts:
        break

    # Get most frequent pair
    best_pair = pair_counts.most_common(1)[0][0]

    # Merge this pair in all words
    for word in word_tokens:
        tokens = word_tokens[word]
        new_tokens = []
        i = 0
        while i < len(tokens):
            if (i < len(tokens) - 1 and
                tokens[i] == best_pair[0] and
                tokens[i + 1] == best_pair[1]):
                # Merge pair
                new_tokens.append(best_pair[0] + best_pair[1])
                i += 2
            else:
                new_tokens.append(tokens[i])
                i += 1
        word_tokens[word] = new_tokens

    # Add merged token to vocabulary
    merged_token = best_pair[0] + best_pair[1]
    self.vocab.append(merged_token)
    self.merges.append(best_pair)
```

: Listing 10.2 — BPE training loop: count all adjacent pairs, merge the most frequent, repeat until the vocabulary reaches its target size. `{#lst-10-tokenization-bpe-train}`

This iterative merging automatically discovers linguistic patterns: common prefixes (“un”, “re”), suffixes (“ing”, “ed”), and frequent words become single tokens. The algorithm requires no linguistic knowledge, learning purely from statistics.

### 18.5.4 Special Tokens

Production tokenizers include special tokens beyond `<UNK>`. Common ones include `<PAD>` for padding sequences to equal length, `<BOS>` (beginning of sequence) and `<EOS>` (end of sequence) for marking bound-

aries, and `<SEP>` for separating multiple text segments. GPT-style models often use `<|endoftext|>` to mark document boundaries.

The choice of special tokens affects the embedding table size. If you reserve 10 special tokens and have a 50,000 token vocabulary, your embedding table has 50,010 rows. Each special token needs learned parameters just like regular tokens.

### 18.5.5 Encoding and Decoding

Encoding converts text to token IDs; decoding reverses the process. For BPE, encoding requires applying learned merge rules in order:

The code in `?@lst-10-tokenization-bpe-encode` makes this concrete.

```
def encode(self, text: str) -> List[int]:
    """Encode text using BPE."""
    # Split text into words
    words = text.split()
    all_tokens = []

    for word in words:
        # Get character-level tokens
        word_tokens = self._get_word_tokens(word)

        # Apply BPE merges
        merged_tokens = self._apply_merges(word_tokens)

        all_tokens.extend(merged_tokens)

    # Convert to IDs
    token_ids = []
    for token in all_tokens:
        token_ids.append(self.token_to_id.get(token, 0)) # 0 = <UNK>

    return token_ids
```

: Listing 10.3 — `BPETokenizer.encode()` splits words, applies the learned merge rules, and looks up integer IDs, falling back to `<UNK>` for unseen tokens. `{#lst-10-tokenization-bpe-encode}`

Decoding is simpler: look up each ID, join the tokens, and clean up markers:

The code in `?@lst-10-tokenization-bpe-decode` makes this concrete.

```
def decode(self, tokens: List[int]) -> str:
    """Decode token IDs back to text."""
    # Convert IDs to tokens
    token_strings = []
    for token_id in tokens:
        token = self.id_to_token.get(token_id, '<UNK>')
        token_strings.append(token)

    # Join and clean up
    text = ''.join(token_strings)
```



```
# Replace end-of-word markers with spaces
text = text.replace('</w>', ' ')

# Clean up extra spaces
text = ' '.join(text.split())

return text
```

: Listing 10.4 — `BPETokenizer.decode()` reverses encoding: look up token strings, join them, and convert `</w>` markers back into whitespace. {#lst-10-tokenization-bpe-decode}

The round-trip text  $\rightarrow$  IDs  $\rightarrow$  text should be lossless for known vocabulary. Unknown tokens degrade gracefully, mapping to `<UNK>` in both directions.

### 18.5.6 Computational Complexity

Character tokenization is fast: encoding is  $O(n)$  where  $n$  is the string length (one dictionary lookup per character), and decoding is also  $O(n)$  (one reverse lookup per ID). The operations are embarrassingly parallel since each character processes independently.

BPE is slower due to merge rule application. Training BPE scales approximately  $O(n^2 \times m)$  where  $n$  is corpus size and  $m$  is the number of merges. Each merge iteration requires counting all pairs across the entire corpus, then updating token sequences. For a 10,000-word corpus learning 5,000 merges, this can take seconds to minutes depending on implementation.

Encoding with trained BPE is  $O(n \times m)$  where  $n$  is text length and  $m$  is the number of merge rules. Each merge rule must scan the token sequence looking for applicable pairs. Production tokenizers optimize this with trie data structures and caching, achieving near-linear time.

Table 18.6 summarises the complexity and speed of each stage.

Table 18.6: Complexity and speed of character and BPE tokenization stages.

| Operation            | Character         | BPE Training            | BPE Encoding           |
|----------------------|-------------------|-------------------------|------------------------|
| <b>Complexity</b>    | $O(n)$            | $O(n^2 \times m)$       | $O(n \times m)$        |
| <b>Typical Speed</b> | 1-5 ms/1K chars   | 100-1000 ms/10K corpus  | 5-20 ms/1K chars       |
| <b>Bottleneck</b>    | Dictionary lookup | Pair frequency counting | Merge rule application |

### 18.5.7 Vocabulary Size Versus Sequence Length

The fundamental trade-off in tokenization creates a spectrum of choices. Small vocabularies (100-500 tokens) produce long sequences because each token represents little information (individual characters or very common subwords). Large vocabularies (50,000+ tokens) produce short sequences because each token represents more information (whole words or meaningful subword units).

Memory and computation scale in opposite directions:

**Embedding table memory** = vocabulary size  $\times$  embedding dimension  $\times$  bytes per parameter  
**Sequence processing cost** = sequence length<sup>2</sup>  $\times$  embedding dimension (for attention)

A character tokenizer (vocab 100, dim 512) needs  $100 \times 512 \times 4 = 200 \text{ KB}$  for embeddings. But a 50-word sentence produces roughly 250 character tokens, requiring  $250^2 = 62,500$  attention computations *per layer*.

A BPE tokenizer (vocab 50,000, dim 512) needs  $50,000 \times 512 \times 4 = 97.7 \text{ MB}$  for embeddings. But the same 50-word sentence collapses to about 75 BPE tokens, requiring  $75^2 = 5,625$  attention computations *per layer*.

A character tokenizer with a vocabulary of 100 and an embedding dimension of 512 requires a negligible `{python} tradeoff_char_embed` for its embedding table. However, a standard 50-word sentence explodes into roughly 250 individual character tokens. Because self-attention complexity scales quadratically with sequence length, this requires  $250^2 = \text{{python} tradeoff_char_attn}$  attention computations *per layer*.

Conversely, a BPE tokenizer expands the vocabulary to 50,000, driving the embedding table footprint up to `{python} tradeoff_bpe_embed`. Yet, that exact same 50-word sentence compresses into a mere 75 BPE subword tokens, slashing the computational burden to just  $75^2 = \text{{python} tradeoff_bpe_attn}$  attention operations per layer.

### **i Systems Implication: The Attention Bottleneck**

The monumental  $O(N^2)$  computational savings in attention (`{python} tradeoff_char_attn` vs. `{python} tradeoff_bpe_attn` ops) absolutely dwarf the static VRAM penalty of hosting a massive embedding table (`{python} tradeoff_char_embed` vs. `{python} tradeoff_bpe_embed`). By condensing text into information-dense subwords, BPE tokenization explicitly trades static GPU memory capacity to aggressively alleviate the catastrophic sequence-length bottleneck in the Transformer's attention mechanism.

This fundamental trade-off dictates the architecture of every production language model: large, static embedding tables are enthusiastically absorbed by VRAM precisely because the resulting abbreviated sequence lengths dramatically accelerate end-to-end training and real-time inference.

Modern language models balance these factors:

Table 18.7 places your implementation in context against production models.

Table 18.7: Vocabulary size and tokenizer strategy in production LLMs.

| Model            | Vocabulary | Strategy      | Sequence Length (typical)     |
|------------------|------------|---------------|-------------------------------|
| <b>GPT-2/3</b>   | 50,257     | BPE           | ~50-200 tokens per sentence   |
| <b>BERT</b>      | 30,522     | WordPiece     | ~40-150 tokens per sentence   |
| <b>T5</b>        | 32,128     | SentencePiece | ~40-180 tokens per sentence   |
| <b>Character</b> | ~100       | Character     | ~250-1000 tokens per sentence |

## 18.6 Production Context

### 18.6.1 Your Implementation vs. Production Tokenizers

Your TinyTorch tokenizers demonstrate the core algorithms, but production tokenizers optimize for speed and scale. The conceptual differences are minimal: the same BPE algorithm, the same vocabulary mappings, the same encode/decode operations. The implementation differences are dramatic.

Table 18.8 places your implementation side by side with the production reference for direct comparison.

Table 18.8: Feature comparison between TinyTorch tokenizers and Hugging Face Tokenizers.

| Feature                | Your Implementation | Hugging Face Tokenizers            |
|------------------------|---------------------|------------------------------------|
| <b>Language</b>        | Pure Python         | Rust (compiled to native code)     |
| <b>Speed</b>           | 1-10 ms/sentence    | 0.01-0.1 ms/sentence (100× faster) |
| <b>Parallelization</b> | Single-threaded     | Multi-threaded with Rayon          |

| Feature                   | Your Implementation | Hugging Face Tokenizers              |
|---------------------------|---------------------|--------------------------------------|
| <b>Vocabulary storage</b> | Python dict         | Trie data structure                  |
| <b>Special features</b>   | Basic encode/decode | Padding, truncation, attention masks |
| <b>Pretrained models</b>  | Train from scratch  | Load from Hugging Face Hub           |

### 18.6.2 Code Comparison

The following comparison shows equivalent tokenization in TinyTorch and Hugging Face. Notice how the high-level API mirrors production tools, making your learning transferable.

## 18.7 Your TinyTorch

```
from tinytorch.core.tokenization import BPETokenizer

# Train tokenizer on corpus
corpus = ["hello world", "machine learning"]
tokenizer = BPETokenizer(vocab_size=1000)
tokenizer.train(corpus)

# Encode text
text = "hello machine"
token_ids = tokenizer.encode(text)

# Decode back to text
decoded = tokenizer.decode(token_ids)
```

## 18.8 Hugging Face

```
from tokenizers import Tokenizer, models, trainers

# Train tokenizer on corpus (same algorithm!)
corpus = ["hello world", "machine learning"]
tokenizer = Tokenizer(models.BPE())
trainer = trainers.BpeTrainer(vocab_size=1000)
tokenizer.train_from_iterator(corpus, trainer)

# Encode text
text = "hello machine"
output = tokenizer.encode(text)
token_ids = output.ids

# Decode back to text
decoded = tokenizer.decode(token_ids)
```

Let's walk through each section to understand the comparison:

- **Lines 1-3 (Imports):** TinyTorch exposes `BPETokenizer` directly from the tokenization module. Hugging Face uses a more modular design with separate `models` and `trainers` for flexibility across algorithms (BPE, WordPiece, Unigram).
- **Lines 5-8 (Training):** Both train on the same corpus using the same BPE algorithm. TinyTorch uses a simpler API with `train()` method. Hugging Face separates model definition from training for composability, but the underlying algorithm is identical.
- **Lines 10-12 (Encoding):** TinyTorch returns a list of integers directly. Hugging Face returns an `Encoding` object with additional metadata (attention masks, offsets, etc.), and you extract the IDs with `.ids` attribute. Same numerical result.
- **Lines 14-15 (Decoding):** Both use `decode()` with the token ID list. Output is identical. The core operation is the same: look up each ID in the vocabulary and join the tokens.

#### 💡 What's Identical

The BPE algorithm, merge rule learning, vocabulary structure, and encode/decode logic. When you debug tokenization issues in production, you'll understand exactly what's happening because you built the same system.

### 18.8.1 Why Tokenization Matters at Scale

The cost of a tokenization choice is invisible on a single example and catastrophic at scale:

- **GPT-3 training:** Processing 300 billion tokens required careful vocabulary selection. Switching to character tokenization would have inflated sequence lengths by 3–4×, multiplying attention cost by 9–16× — adding millions of dollars to a single training run.
- **Embedding table memory:** A 50,000-token vocabulary with 12,288-dimensional embeddings (GPT-3) is  $50,000 \times 12,288 \times 4$  bytes = **2.29 GB** for the embedding layer alone. That is only ~0.35% of GPT-3's 175 B parameters, yet it is the layer that touches every input on every forward pass.
- **Real-time inference:** Chatbots must tokenize user input in milliseconds. Python tokenizers take 5–20 ms per sentence; Rust tokenizers take 0.05–0.2 ms. At 1 million requests per day, the difference is roughly 5 hours of CPU time — every day.

#### 💡 Check Your Understanding — Tokenization

Before moving on, verify you can articulate each of the following:

- The memory/sequence-length trade-off: BPE's 50K-vocab embedding table (tens of MB) versus character-level's  $O(n^2)$  attention bloat from long sequences.
- Why BPE training is  $O(n^2 \times m)$  while trained BPE *encoding* is near-linear, and why production tokenizers run in Rust rather than Python.
- How `<UNK>` fallback and end-of-word markers make round-trip encode/decode lossy for unseen characters but deterministic for known vocabulary.
- Why the vocabulary you pick upstream caps every downstream choice: embedding-table parameters, context window in tokens, and attention cost per step.

If any of these feels fuzzy, revisit the *Core Concepts* section (especially *Vocabulary Size Versus Sequence Length*) before moving on.

## 18.9 Self-Check Questions

Test yourself with these systems thinking questions. They're designed to build intuition for the performance characteristics you'll encounter in production NLP systems.

### Q1: Vocabulary Memory Calculation

You train a BPE tokenizer with `vocab_size=30,000` for a production model. If using 768-dimensional embeddings with float32 precision, how much memory does the embedding table require?

#### 💡 Answer

$30,000 \times 768 \times 4 \text{ bytes} = 92,160,000 \text{ bytes} \approx 87.89 \text{ MB}$

Breakdown:

- Vocabulary size: 30,000 tokens
- Embedding dimension: 768 (BERT-base size)
- Float32: 4 bytes per parameter
- Total parameters:  $30,000 \times 768 = 23,040,000$
- Memory:  $23.04\text{M} \times 4 = 87.89 \text{ MB}$

This is why vocabulary size matters. Doubling to 60K vocab doubles embedding memory to ~176 MB — and you pay it once at load time, every time, on every device.

### Q2: Sequence Length Trade-offs

A sentence contains 200 characters. With character tokenization it produces 200 tokens. With BPE it produces 50 tokens (4:1 compression). If processing batch size 32 with attention:

- How many attention computations for character tokenization per batch?
- How many for BPE tokenization per batch?

#### 💡 Answer

##### Character tokenization:

- Sequence length: 200 tokens
- Attention per sequence:  $200^2 = 40,000$  operations
- Batch size: 32
- Total:  $32 \times 40,000 = 1,280,000$  attention operations

##### BPE tokenization:

- Sequence length: 50 tokens (200 chars  $\div$  4)
- Attention per sequence:  $50^2 = 2,500$  operations
- Batch size: 32
- Total:  $32 \times 2,500 = 80,000$  attention operations

BPE is **16 $\times$  faster** at attention. This is why modern models use subword tokenization despite the larger embedding table.

### Q3: Unknown Token Handling

Your BPE tokenizer encounters the word “supercalifragilistic” (not in training corpus). Character tokenizer maps it to 20 known tokens. BPE tokenizer decomposes it into subwords like ['super', 'cal', 'ifr', 'ag', 'il', 'istic'] (6 tokens). Which is better?

 **Answer****BPE is better for production:**

- **Efficiency:** 6 tokens vs 20 tokens — a  $3.3\times$  shorter sequence
- **Semantics:** Subwords like “super” and “istic” carry meaning; individual characters don’t
- **Generalization:** Model learns that the “super” prefix modifies meaning (superman, supermarket)
- **Compute:** 36 attention computations vs 400 ( $\sim 11\times$  fewer)

**Character tokenization advantages:**

- **Perfect coverage:** Never emits `<UNK>`, always reconstructs the original text
- **Simplicity:** No merge rules, no training step

For rare or out-of-vocabulary words, BPE’s subword decomposition wins on both efficiency and semantics, which is why GPT, BERT, and T5 all use variants of subword tokenization.

**Q4: Compression Ratio Analysis**

You analyze two tokenizers on a 10,000 character corpus:

- Character tokenizer: 10,000 tokens
- BPE tokenizer: 2,500 tokens

What’s the compression ratio, and what does it tell you about efficiency?

 **Answer**

**Compression ratio:**  $10,000 \div 2,500 = 4.0$

Each BPE token represents an average of 4 characters.

**Efficiency implications:**

- **Sequence processing:**  $4.0\times$  shorter sequences  $\rightarrow 16\times$  faster attention (quadratic scaling)
- **Context window:** With max length 512, the character tokenizer fits 512 chars ( $\sim 100$  words); BPE fits 2,048 chars ( $\sim 400$  words)
- **Information density:** Each BPE token carries more semantic information (subword vs character)

**Trade-off:** BPE vocabulary is  $\sim 100\times$  larger (10K tokens vs 100), pushing embedding memory from  $\sim 200$  KB to  $\sim 20$  MB. The trade favors BPE heavily once you stack multiple transformer layers and attention dominates the cost.

**Q5: Training Corpus Size Impact**

Training BPE on 1,000 words takes 100ms. How long will 10,000 words take? What about 100,000 words?

 **Answer**

BPE training scales approximately  $O(n^2)$  where  $n$  is corpus size (each merge re-counts every pair across the corpus).

- **1,000 words:** 100 ms (baseline)
- **10,000 words:**  $\sim 10,000$  ms = 10 seconds ( $100\times$  longer, from  $10^2$  scaling)
- **100,000 words:**  $\sim 1,000,000$  ms = 1,000 seconds  $\approx 16.7$  minutes ( $10,000\times$  longer)

**Production strategies to handle this:**

- Sample a representative subset ( $\sim 50K$ - $100K$  sentences is usually enough)
- Use incremental/online BPE that doesn’t recount all pairs each iteration
- Parallelize pair counting across corpus chunks

- Cache frequent pair statistics
- Use optimized implementations (Rust, C++) that are 100–1,000× faster

Note: encoding with a *trained* BPE tokenizer is fast (near-linear). Only the training phase is expensive.

## 18.10 Key Takeaways

- **Tokens, not text, are the model’s input:** every upstream vocabulary decision silently sets sequence length, embedding-table size, and per-step attention cost.
- **Vocabulary size is a memory/compute trade:** small vocabularies save embedding memory but blow up attention (quadratic in sequence length); large vocabularies pay MBs of embeddings to keep sequences short.
- **BPE learns subwords from statistics alone:** iteratively merging the most frequent adjacent pair yields linguistically meaningful units (prefixes, suffixes, whole words) without any language-specific rules.
- **Training is expensive; encoding is cheap:** BPE training is  $\sim O(n^2 \times m)$ , but once merges are learned, encoding is near-linear — which is why production tokenizers ship as pre-trained artifacts in Rust/C++.

**Coming next:** Module 11 turns the integer token IDs you produced here into learnable dense vectors — the first place gradients actually touch language.

## 18.11 Further Reading

For students who want to understand the academic foundations and production implementations of tokenization:

### 18.11.1 Seminal Papers

- **Neural Machine Translation of Rare Words with Subword Units** - Sennrich et al. (2016). The original BPE paper that introduced subword tokenization for neural machine translation. Shows how BPE handles rare words through intelligent subword decomposition, achieving superior translation quality.
  - **Systems Implication:** By mathematically compressing raw text into information-dense subword tokens, this algorithm radically reduced average sequence lengths. This directly mitigated the crippling  $O(N^2)$  memory and compute bottlenecks inherent to sequence-processing attention layers, enabling the modern LLM era. [arXiv:1508.07909](#)
- **SentencePiece: A simple and language independent approach to subword tokenization** - Kudo & Richardson (2018). Extends BPE with a language-agnostic, raw-text-to-token pipeline that elegantly sidesteps brittle regex pre-tokenization rules. Used extensively in T5, ALBERT, and LLaMA.
  - **Systems Implication:** SentencePiece cleanly eliminated complex, language-specific text pre-processing heuristics from the dataloader. By absorbing the entire pipeline into a unified, heavily optimized C++ binary, it prevented host CPUs from bottlenecking the GPU during massive-scale text ingestion. [arXiv:1808.06226](#)
- **BERT: Pre-training of Deep Bidirectional Transformers** - Devlin et al. (2018). While primarily focused on bidirectional transformer encoders, this paper formally introduced WordPiece tokenization to the mass-market NLP ecosystem.
  - **Systems Implication:** By aggressively fixing the vocabulary size at 30,522 tokens, BERT strictly constrained the memory footprint of the final, massive softmax projection layer. This calculated engineering decision ensured the enormous output embedding matrix remained tightly bounded within the strict VRAM limits of 2018-era datacenter GPUs. [arXiv:1810.04805](#)

### 18.11.2 Additional Resources

- **Library:** [Hugging Face Tokenizers](#) - Production Rust implementation with Python bindings. Explore the source to see optimized BPE.
- **Tutorial:** “Byte Pair Encoding Tokenization” - [Hugging Face Course](#) - Interactive tutorial showing BPE in action with visualizations
- **Textbook:** “[Speech and Language Processing](#)” by Jurafsky & Martin - Chapter 2 covers tokenization, including Unicode handling and language-specific issues

## 18.12 What’s Next

### Coming Up: Module 11 — Embeddings

You now have integer token IDs. Integers, however, are *opaque*: ID 142 and ID 143 are no more related than ID 142 and ID 9,001. The next module turns those IDs into **learnable dense vectors** so that “king” and “queen” can sit close together in space, and so that gradients can actually flow through your model. The vocabulary you fixed in this module sets the number of rows in that embedding table — and therefore the bulk of the parameters at the input of every transformer you’ll build.

#### Preview — how your tokenizer feeds the rest of the stack:

Table 18.9 traces how this module is reused by later parts of the curriculum.

Table 18.9: How tokenization feeds into embeddings, attention, and transformers.

| Module                  | What It Does                    | Your Tokenization In Action                                  |
|-------------------------|---------------------------------|--------------------------------------------------------------|
| <b>11: Embeddings</b>   | Learnable lookup tables         | <code>embedding = Embedding(vocab_size=1000, dim=128)</code> |
| <b>12: Attention</b>    | Sequence-to-sequence processing | Token sequences attend to each other                         |
| <b>13: Transformers</b> | Complete language models        | Full pipeline: tokenize → embed → attend → predict           |

## 18.13 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.

Transformers\*\* | Complete language models | Full pipeline: tokenize → embed → attend → predict |



## 🔥 Chapter 19

# Module 11: Embeddings

Embeddings are the largest single tensor in most language models (GPT-3's 2.3 GB) but every forward pass only touches a handful of rows. That makes them a row-sparse scatter/gather problem, not a matmul, and it puts them squarely in the memory-bandwidth regime of the roofline. This module builds the lookup table, the positional encodings, and the sparse-gradient backward pass that every transformer input layer depends on.

### Module Info

**ARCHITECTURE TIER** | Difficulty: ●●○○ | Time: 3-5 hours | Prerequisites: 01-08, 10

**Prerequisites:** Modules 01-08 and 10 means you should understand:

- Tensor operations (shape manipulation, matrix operations, broadcasting)
- Training fundamentals (forward/backward, optimization)
- Tokenization (converting text to token IDs, vocabularies)

If you can explain how a tokenizer converts “hello” to token IDs and how to multiply matrices, you’re ready.

## 19.1 Overview

Neural networks operate on vectors. Language is made of tokens. Embeddings are how the two meet: a learnable lookup table that turns each integer token ID into a dense vector, so the rest of the network can do calculus on it.

Your tokenizer from Module 10 produces IDs like `[42, 7, 15]`. By the end of this module you have built the layer that turns those IDs into geometry — the same layer sitting at the input of every transformer from BERT to GPT-4 — and the positional encodings that tell the network *where* in the sequence each token lives.

## 19.2 Learning Objectives

### By completing this module, you will:

- **Implement** embedding layers that convert token IDs to dense vectors through efficient table lookup
- **Master** positional encoding strategies including learned and sinusoidal approaches
- **Understand** memory scaling for embedding tables and the trade-offs between vocabulary size and embedding dimension
- **Connect** your implementation to production transformer architectures used in GPT and BERT




Figure 19.1: **TinyTorch Embedding System:** Mapping tokens and positions to learned continuous representations.

## 19.3 What You'll Build

### Implementation roadmap:

Table 19.1 lays out the implementation in order, one part at a time.

Table 19.1: **Implementation roadmap for the embedding classes.**

| Part | What You'll Implement                       | Key Concept                            |
|------|---------------------------------------------|----------------------------------------|
| 1    | <code>Embedding</code> class                | Token ID to vector lookup via indexing |
| 2    | <code>PositionalEncoding</code> class       | Learnable position embeddings          |
| 3    | <code>create_sinusoidal_embeddings()</code> | Mathematical position encoding         |
| 4    | <code>EmbeddingLayer</code> class           | Complete token + position system       |

### The pattern you'll enable:

```
# Converting tokens to position-aware dense vectors
embed_layer = EmbeddingLayer(vocab_size=50000, embed_dim=512)
tokens = Tensor([[1, 42, 7]]) # Token IDs from tokenizer
embeddings = embed_layer(tokens) # (1, 3, 512) dense vectors ready for attention
```

### 19.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Attention mechanisms (that's Module 12: Attention)
- Full transformer architectures (that's Module 13: Transformers)

- Word2Vec or GloVe pretrained embeddings (you're building learnable embeddings)
- Subword embedding composition (PyTorch handles this at the tokenization level)

**You are building the foundation for sequence models.** Context-aware representations come next.

## 19.4 API Reference

This section documents the embedding components you'll build. Use this as your reference while implementing and debugging.

### 19.4.1 Embedding Class

```
Embedding(vocab_size, embed_dim)
```

Learnable embedding layer that maps token indices to dense vectors through table lookup.

**Constructor Parameters:** - `vocab_size` (int): Size of vocabulary (number of unique tokens) - `embed_dim` (int): Dimension of embedding vectors

**Core Methods:**

Table 19.2 lists the methods on this class.

Table 19.2: Core methods on the Embedding class.

| Method                  | Signature                                          | Description                           |
|-------------------------|----------------------------------------------------|---------------------------------------|
| <code>forward</code>    | <code>forward(indices: Tensor) -&gt; Tensor</code> | Lookup embeddings for token indices   |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code>       | Return weight matrix for optimization |

**Properties:** - `weight`: Tensor of shape `(vocab_size, embed_dim)` containing learnable embeddings

### 19.4.2 PositionalEncoding Class

```
PositionalEncoding(max_seq_len, embed_dim)
```

Learnable positional encoding that adds trainable position-specific vectors to embeddings.

**Constructor Parameters:** - `max_seq_len` (int): Maximum sequence length to support - `embed_dim` (int): Embedding dimension (must match token embeddings)

**Core Methods:**

Table 19.3 lists the methods on this class.

Table 19.3: Core methods on the PositionalEncoding class.

| Method                  | Signature                                    | Description                                  |
|-------------------------|----------------------------------------------|----------------------------------------------|
| <code>forward</code>    | <code>forward(x: Tensor) -&gt; Tensor</code> | Add positional encodings to input embeddings |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code> | Return position embedding matrix             |

### 19.4.3 Sinusoidal Embeddings Function

```
create_sinusoidal_embeddings(max_seq_len, embed_dim) -> Tensor
```

Creates fixed sinusoidal positional encodings using trigonometric functions. No parameters to learn.

#### Mathematical Formula:

$$PE(pos, 2i) = \sin(pos / 10000^{(2i/embed\_dim)})$$

$$PE(pos, 2i+1) = \cos(pos / 10000^{(2i/embed\_dim)})$$

### 19.4.4 EmbeddingLayer Class

```
EmbeddingLayer(vocab_size, embed_dim, max_seq_len=512,
               pos_encoding='learned', scale_embeddings=False)
```

Complete embedding system combining token embeddings and positional encoding.

**Constructor Parameters:** - `vocab_size` (int): Size of vocabulary - `embed_dim` (int): Embedding dimension - `max_seq_len` (int): Maximum sequence length for positional encoding - `pos_encoding` (str): Type of positional encoding ('learned', 'sinusoidal', or None) - `scale_embeddings` (bool): Whether to scale by  $\sqrt{\text{embed\_dim}}$  (transformer convention)

#### Core Methods:

Table 19.4 lists the methods on this class.

Table 19.4: Core methods on the combined token+position Embedding wrapper.

| Method                  | Signature                                         | Description                 |
|-------------------------|---------------------------------------------------|-----------------------------|
| <code>forward</code>    | <code>forward(tokens: Tensor) -&gt; Tensor</code> | Complete embedding pipeline |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code>      | All trainable parameters    |

## 19.5 Core Concepts

Four ideas carry the rest of this chapter: lookup-as-matmul, sparse gradient updates, the learned-vs-fixed positional tradeoff, and the memory cost of embedding dimension. Every modern language model is built from them.

### 19.5.1 From Indices to Vectors

Neural networks consume continuous vectors; language arrives as discrete IDs. After the tokenizer turns “the cat sat” into `[1, 42, 7]`, you need a way to give those integers geometry.

The embedding layer is that bridge: a learnable matrix where row `i` is the vector for token `i`. For a 50,000-token vocabulary with 512-dimensional embeddings, the table is shape `(50000, 512)`, initialized small. Training nudges similar words toward similar rows — semantic relationships emerge as geometric proximity.

The implementation is one line of NumPy:

The code in `?@lst-11-embedding-forward` makes this concrete.

```
def forward(self, indices: Tensor) -> Tensor:
    """Forward pass: lookup embeddings for given indices."""
```

```

# Validate indices are in range
if np.any(indices.data >= self.vocab_size) or np.any(indices.data < 0):
    raise ValueError(
        f"Index out of range. Expected 0 <= indices < {self.vocab_size}, "
        f"got min={np.min(indices.data)}, max={np.max(indices.data)}"
    )

# Perform embedding lookup using advanced indexing
# This is equivalent to one-hot multiplication but much more efficient
embedded = self.weight.data[indices.data.astype(int)]

# Create result tensor with gradient tracking
result = Tensor(embedded, requires_grad=self.weight.requires_grad)

if result.requires_grad:
    result._grad_fn = EmbeddingBackward(self.weight, indices)

return result

```

: **Listing 11.1 — Embedding forward pass.** NumPy fancy indexing replaces a one-hot matmul, making lookup  $O(1)$  per token. {#lst-11-embedding-forward}

The work happens in `self.weight.data[indices.data.astype(int)]`. NumPy’s fancy indexing pulls rows 1, 42, and 7 in a single operation, broadcasting cleanly over batched inputs of any shape. It is mathematically equivalent to multiplying a one-hot matrix against the weight table — but it is orders of magnitude faster and allocates no intermediate tensor.

### 19.5.2 Embedding Table Mechanics

The embedding table is a learnable parameter matrix initialized with small random values. For vocabulary size  $V$  and embedding dimension  $D$ , the table has shape  $(V, D)$ . Each row represents one token’s learned representation.

Initialization matters for training stability. The implementation uses Xavier initialization:

```

# Xavier initialization for better gradient flow
limit = math.sqrt(6.0 / (vocab_size + embed_dim))
self.weight = Tensor(
    np.random.uniform(-limit, limit, (vocab_size, embed_dim)),
    requires_grad=True
)

```

This initialization scale ensures gradients neither explode nor vanish at the start of training. The limit is computed from both vocabulary size and embedding dimension, balancing the fan-in and fan-out of the embedding layer.

During training, gradients flow back through the lookup operation to update only the embeddings that were accessed. If your batch contains tokens `[5, 10, 10, 5]`, only rows 5 and 10 of the embedding table receive gradient updates. This sparse gradient pattern is handled by the `EmbeddingBackward` gradient function, making embedding updates extremely efficient even for vocabularies with millions of tokens.

**i Systems Implication: Sparse Gradient Updates and Embedding Memory Layout**

The embedding table is the largest *single* tensor in most language models — GPT-3's token embeddings alone consume 2.3 GB — yet a forward pass only touches the handful of rows corresponding to tokens in the current batch. Treating the embedding step as a **row-sparse scatter/gather** (rather than a dense matmul against a one-hot matrix) is what keeps training tractable: you never materialize the  $(\text{batch} \times \text{seq}, \text{vocab})$  one-hot tensor, you never multiply by it, and the backward pass only writes gradients into the rows that were read. Production frameworks go further — they store the embedding matrix in row-major layout so each lookup is one contiguous HBM transaction per token, use sparse gradient tensors in the optimizer (Adam state for only the accessed rows), and in distributed training shard the table across devices so no single GPU has to hold all 2.3 GB. Every one of these optimizations is a direct consequence of the lookup-as-sparse-scatter pattern you just implemented.

**19.5.3 Learned vs Fixed Embeddings**

Positional information can be added to token embeddings in two fundamentally different ways: learned position embeddings and fixed sinusoidal encodings.

**Learned positional encoding** treats each sequence position as a unique, trainable parameter, just like token embeddings. For a maximum sequence length  $M$  and embedding dimension  $D$ , you construct an auxiliary embedding table of shape  $(M, D)$ :

```
# Initialize position embedding matrix
# Smaller initialization than token embeddings since these are additive
limit = math.sqrt(2.0 / embed_dim)
self.position_embeddings = Tensor(
    np.random.uniform(-limit, limit, (max_seq_len, embed_dim)),
    requires_grad=True
)
```

During forward pass, you slice position embeddings and add them to token embeddings: The code in `?lst-11-positional-forward` makes this concrete.

```
def forward(self, x: Tensor) -> Tensor:
    """Add positional encodings to input embeddings."""
    batch_size, seq_len, embed_dim = x.shape

    # Slice position embeddings for this sequence length
    # Tensor slicing preserves gradient flow (from Module 01's __getitem__)
    pos_embeddings = self.position_embeddings[:seq_len]

    # Reshape to add batch dimension: (1, seq_len, embed_dim)
    pos_data = pos_embeddings.data[np.newaxis, :, :]
    pos_embeddings_batched = Tensor(pos_data,
    requires_grad=pos_embeddings.requires_grad)

    # Copy gradient function to preserve backward connection
    if hasattr(pos_embeddings, '_grad_fn') and pos_embeddings._grad_fn is not None:
        pos_embeddings_batched._grad_fn = pos_embeddings._grad_fn
```

```
# Add positional information - gradients flow through both x and pos_embeddings!
result = x + pos_embeddings_batched
return result
```

: **Listing 11.2 — PositionalEncoding forward pass.** Slices position embeddings to current sequence length and adds them to token embeddings while preserving gradient flow. {#lst-11-positional-forward}

The slicing operation `self.position_embeddings[:seq_len]` preserves gradient tracking because TinyTorch’s Tensor `__getitem__` (from Module 01) maintains the connection to the original parameter. This allows backpropagation to update only the position embeddings actually used in the forward pass.

The advantage is flexibility: the model can learn task-specific positional patterns. The disadvantage is memory cost and a hard maximum sequence length.

**Fixed sinusoidal encoding** uses mathematical patterns requiring no parameters. The formula creates unique position signatures using sine and cosine at different frequencies:

The code in `?@lst-11-sinusoidal-pe` makes this concrete.

```
# Create position indices [0, 1, 2, ..., max_seq_len-1]
position = np.arange(max_seq_len, dtype=np.float32)[: , np.newaxis]

# Create dimension indices for calculating frequencies
div_term = np.exp(
    np.arange(0, embed_dim, 2, dtype=np.float32) *
    -(math.log(10000.0) / embed_dim)
)

# Initialize the positional encoding matrix
pe = np.zeros((max_seq_len, embed_dim), dtype=np.float32)

# Apply sine to even indices, cosine to odd indices
pe[:, 0::2] = np.sin(position * div_term)
pe[:, 1::2] = np.cos(position * div_term)
```

: **Listing 11.3 — Sinusoidal positional encoding construction.** Builds a fixed position matrix from trigonometric functions at geometrically spaced frequencies. {#lst-11-sinusoidal-pe}

This creates a unique vector for each position where low-index dimensions oscillate rapidly (high frequency) and high-index dimensions change slowly (low frequency). The pattern allows the model to learn relative positions through dot products, and crucially, can extrapolate to sequences longer than seen during training.

### 19.5.4 Positional Encodings

Token embeddings are a bag of vectors — they carry no notion of order. Without positional information, “cat sat on mat” and “mat on sat cat” produce the same set of vectors and a downstream attention layer cannot tell them apart.

Positional encodings fix this by adding a position-specific vector to each token’s embedding. Token 42 at position 0 lands somewhere different in space than token 42 at position 5, and the network finally has access to *where*, not just *what*.

The Transformer paper introduced sinusoidal positional encoding with a clever mathematical structure. For position `pos` and dimension `i`:

```
PE(pos, 2i)    = sin(pos / 10000^(2i/embed_dim)) # Even dimensions
PE(pos, 2i+1)  = cos(pos / 10000^(2i/embed_dim)) # Odd dimensions
```

The 10000 base creates different wavelengths across dimensions. Dimension 0 oscillates rapidly (frequency  $\approx 1$ ), while dimension 510 changes extremely slowly (frequency  $\approx 1/10000$ ). This multiscale structure gives each position a unique “fingerprint” and enables the model to learn relative position through simple vector arithmetic.

At position 0, all sine terms equal 0 and all cosine terms equal 1:  $[0, 1, 0, 1, 0, 1, \dots]$ . At position 1, the pattern shifts based on each dimension’s frequency. The combination of many frequencies creates distinct encodings where nearby positions have similar (but not identical) vectors, providing smooth positional gradients.

The trigonometric identity enables learning relative positions:  $PE(pos+k)$  can be expressed as a linear function of  $PE(pos)$  using sine and cosine addition formulas. This allows attention mechanisms to implicitly learn positional offsets (like “the token 3 positions ahead”) through learned weights on the position encodings, without the model needing separate relative position parameters.

### 19.5.5 Embedding Dimension Trade-offs

The embedding dimension  $D$  controls the capacity of your learned representations. Larger  $D$  provides more expressiveness but costs memory and compute. The choice involves several interacting factors.

**Memory scaling:** Embedding tables scale as  $vocab\_size \times embed\_dim \times 4 \text{ bytes (float32)}$ . A 50,000-token vocabulary with 512-dimensional embeddings costs **98 MB**. Double the dimension to 1024 and the cost doubles to **195 MB**. For large vocabularies, the embedding table often dominates total model memory — GPT-3’s 50,257-token vocabulary with 12,288-dimensional embeddings consumes **2.3 GB** for token embeddings alone.

**Semantic capacity:** Higher dimensions allow finer-grained semantic distinctions. With 64 dimensions, you might capture basic categories (animals, actions, objects). With 512 dimensions, you can encode subtle relationships (synonyms, antonyms, part-of-speech, contextual variations). With 1024+ dimensions, you have capacity for highly nuanced semantic features discovered through training.

**Computational cost:** Every attention head in transformers performs operations over the embedding dimension. Memory bandwidth becomes the bottleneck: transferring embedding vectors from RAM to cache dominates the time to process them. Larger embeddings mean more memory traffic per token, reducing throughput.

#### Typical scales in production:

Table 19.5 places your implementation in context against production models.

Table 19.5: Vocabulary size, embedding dimension, and memory cost across production LLMs.

| Model             | Vocabulary | Embed Dim | Embedding Memory |
|-------------------|------------|-----------|------------------|
| Small BERT        | 30,000     | 768       | 88 MB            |
| GPT-2             | 50,257     | 1,024     | 196 MB           |
| GPT-3             | 50,257     | 12,288    | 2,356 MB         |
| Large Transformer | 100,000    | 1,024     | 391 MB           |

The embedding dimension typically matches the model’s hidden dimension since embeddings feed directly into the first transformer layer. You rarely see models with embedding dimension different from hidden dimension (though it’s technically possible with a projection layer).



## 19.6 Common Errors

These are the errors you'll encounter most often when working with embeddings. Understanding why they happen will save you hours of debugging.

### 19.6.1 Index Out of Range

**Error:** `ValueError: Index out of range. Expected 0 <= indices < 50000, got max=50001`

Embedding lookup expects token IDs in the range `[0, vocab_size-1]`. If your tokenizer produces an ID of 50001 but your embedding layer has `vocab_size=50000`, the lookup fails.

**Cause:** Mismatch between tokenizer vocabulary size and embedding layer vocabulary size. This often happens when you train a tokenizer separately and forget to sync the vocabulary size when creating the embedding layer.

**Fix:** Ensure `embed_layer.vocab_size` matches your tokenizer's vocabulary size exactly:

```
tokenizer = Tokenizer(vocab_size=50000)
embed = Embedding(vocab_size=tokenizer.vocab_size, embed_dim=512)
```

### 19.6.2 Sequence Length Exceeds Maximum

**Error:** `ValueError: Sequence length 1024 exceeds maximum 512`

Learned positional encodings have a fixed maximum sequence length set during initialization. If you try to process a sequence longer than this maximum, the forward pass fails because there are no position embeddings for those positions.

**Cause:** Input sequences longer than `max_seq_len` parameter used when creating the positional encoding layer.

**Fix:** Either increase `max_seq_len` during initialization, truncate your sequences, or use sinusoidal positional encoding which can handle arbitrary lengths:

```
# Option 1: Increase max_seq_len
pos_enc = PositionalEncoding(max_seq_len=2048, embed_dim=512)

# Option 2: Use sinusoidal (no length limit)
embed_layer = EmbeddingLayer(vocab_size=50000, embed_dim=512,
                             pos_encoding='sinusoidal')
```

### 19.6.3 Embedding Dimension Mismatch

**Error:** `ValueError: Embedding dimension mismatch: expected 512, got 768`

When adding positional encodings to token embeddings, the dimensions must match exactly. If your token embeddings are 512-dimensional but your positional encoding expects 768-dimensional inputs, element-wise addition fails.

**Cause:** Creating embedding components with inconsistent `embed_dim` values.

**Fix:** Use the same `embed_dim` for all embedding components:

```
embed_dim = 512
token_embed = Embedding(vocab_size=50000, embed_dim=embed_dim)
pos_enc = PositionalEncoding(max_seq_len=512, embed_dim=embed_dim)
```

### 19.6.4 Shape Errors with Batching

**Error:** `ValueError: Expected 3D input (batch, seq, embed), got shape (128, 512)`

Positional encoding expects 3D tensors with batch dimension. If you pass a 2D tensor (sequence, embedding), the forward pass fails.

**Cause:** Forgetting to add batch dimension when processing single sequences, or using raw embedding output without reshaping.

**Fix:** Ensure inputs have batch dimension, even for single sequences:

```
# Wrong: 2D input
tokens = Tensor([1, 2, 3])
embeddings = embed(tokens) # Shape: (3, 512) - missing batch dim

# Right: 3D input
tokens = Tensor([[1, 2, 3]]) # Added batch dimension
embeddings = embed(tokens) # Shape: (1, 3, 512) - correct
```

## 19.7 Production Context

### 19.7.1 Your Implementation vs. PyTorch

Your TinyTorch embedding system and PyTorch's `torch.nn.Embedding` share the same conceptual design and API patterns. The differences are in scale, optimization, and device support.

Table 19.6 places your implementation side by side with the production reference for direct comparison.

Table 19.6: Feature comparison between TinyTorch embeddings and `torch.nn.Embedding`.

| Feature                    | Your Implementation  | PyTorch                        |
|----------------------------|----------------------|--------------------------------|
| <b>Backend</b>             | NumPy (CPU only)     | C++/CUDA (CPU/GPU)             |
| <b>Lookup Speed</b>        | 1x (baseline)        | 10-100x faster on GPU          |
| <b>Max Vocabulary</b>      | Limited by RAM       | Billions (with techniques)     |
| <b>Positional Encoding</b> | Learned + sinusoidal | Must implement yourself*       |
| <b>Sparse Gradients</b>    | Via custom backward  | Native sparse gradient support |
| <b>Memory Optimization</b> | Standard             | Quantization, pruning, sharing |

\*PyTorch provides building blocks but you implement positional encoding patterns yourself (as you did here)

### 19.7.2 Code Comparison

The following comparison shows equivalent embedding operations in TinyTorch and PyTorch. Notice how the APIs mirror each other closely.

## 19.8 Your TinyTorch

```

from tinytorch.core.embeddings import Embedding, EmbeddingLayer

# Create embedding layer
embed = Embedding(vocab_size=50000, embed_dim=512)

# Token lookup
tokens = Tensor([[1, 42, 7, 99]])
embeddings = embed(tokens)  # (1, 4, 512)

# Complete system with position encoding
embed_layer = EmbeddingLayer(
    vocab_size=50000,
    embed_dim=512,
    max_seq_len=2048,
    pos_encoding='learned'
)
position_aware = embed_layer(tokens)

```

## 19.9 PyTorch

```

import torch
import torch.nn as nn

# Create embedding layer
embed = nn.Embedding(num_embeddings=50000, embedding_dim=512)

# Token lookup
tokens = torch.tensor([[1, 42, 7, 99]])
embeddings = embed(tokens)  # (1, 4, 512)

# Complete system (you implement positional encoding yourself)
class EmbeddingWithPosition(nn.Module):
    def __init__(self, vocab_size, embed_dim, max_seq_len):
        super().__init__()
        self.token_embed = nn.Embedding(vocab_size, embed_dim)
        self.pos_embed = nn.Embedding(max_seq_len, embed_dim)

    def forward(self, tokens):
        seq_len = tokens.shape[1]
        positions = torch.arange(seq_len).unsqueeze(0)
        return self.token_embed(tokens) + self.pos_embed(positions)

embed_layer = EmbeddingWithPosition(50000, 512, 2048)
position_aware = embed_layer(tokens)

```

Let's walk through each section to understand the comparison:

- **Line 1-2 (Import):** Both frameworks provide dedicated embedding modules. TinyTorch packages everything in `embeddings`; PyTorch uses `nn.Embedding` as the base class.
- **Line 4-5 (Embedding Creation):** Your `Embedding` class closely mirrors PyTorch's `nn.Embedding`. The parameter names differ (`vocab_size` vs `num_embeddings`) but the concept is identical.
- **Line 7-9 (Token Lookup):** Both use identical calling patterns. The embedding layer acts as a function, taking token IDs and returning dense vectors. Shape semantics are identical.
- **Line 11-20 (Complete System):** Your `EmbeddingLayer` provides a complete system in one class. In PyTorch, you implement this pattern yourself by composing `nn.Embedding` layers for tokens and positions. The HuggingFace Transformers library implements this exact pattern for BERT, GPT, and other models.
- **Line 22-24 (Forward Pass):** Both systems add token and position embeddings element-wise. Your implementation handles this internally; PyTorch requires you to manage position indices explicitly.

#### 💡 What's Identical

Embedding lookup semantics, gradient flow patterns, and the addition of positional information. When you debug PyTorch transformer models, you'll recognize these exact patterns because you built them yourself.

### 19.9.1 Why Embeddings Matter at Scale

The numbers behind a modern language model make the design pressures concrete:

- **GPT-3 embeddings:**  $50,257 \text{ tokens} \times 12,288 \text{ dimensions} = 618 \text{ million parameters} = 2.3 \text{ GB}$  for token embeddings alone (positional encodings extra).
- **Lookup throughput:** A batch of 32 sequences  $\times$  2048 tokens demands **65,536 embedding lookups** per step. At 1,000 steps/sec, that is roughly 65 million lookups per second.
- **Memory bandwidth:** Each lookup pulls 2–4 KB from RAM to cache ( $512\text{--}1024 \text{ floats} \times 4 \text{ bytes}$ ). At scale, bandwidth — not arithmetic — is the bottleneck.
- **Gradient sparsity:** A batch touches only a tiny slice of the 50,257-row table. Efficient trainers exploit this and update gradients only for the rows that were accessed.

Embeddings take only **10–15% of training time** but consume **30–40% of model memory** at large vocabularies. They are cheap to *run* and expensive to *store* — exactly the opposite profile of attention.

#### 💡 Check Your Understanding — Embeddings

Before moving on, verify you can articulate each of the following:

- Why the embedding table is a **lookup** (sparse row scatter/gather), not a matmul against a one-hot matrix, and how this reshapes training memory: only the touched rows receive gradient updates.
- How memory scales: embedding storage is `vocab  $\times$  dim  $\times$  bytes` (linear in both), while attention cost downstream is quadratic in *sequence length* — two independent budgets that fight for the same VRAM.
- The learned-vs-sinusoidal positional encoding trade: learned PE is flexible but has a hard `max_seq_len` ceiling; sinusoidal PE is parameter-free and extrapolates, enabling variable-length inference.
- Why GPT-3's 2.3 GB embedding table is only ~0.35% of its parameters but sits on the **critical path of every forward pass**, making embedding lookup a memory-bandwidth problem rather than a FLOP problem.

If any of these feels fuzzy, revisit the *Core Concepts* section (especially *Embedding Table Mechanics* and *Embedding Dimension Trade-offs*) before moving on.

## 19.10 Self-Check Questions

Test yourself with these systems thinking questions. They're designed to build intuition for the performance and memory characteristics you'll encounter in production.

### Q1: Memory Calculation

An embedding layer has `vocab_size=50000` and `embed_dim=512`. How much memory does the embedding table use (in MB)?

#### 💡 Answer

$50,000 \times 512 \times 4 \text{ bytes} = 102,400,000 \text{ bytes} = 97.7 \text{ MB}$

Calculation breakdown:

- Parameters:  $50,000 \times 512 = 25,600,000$
- Memory:  $25,600,000 \times 4 \text{ bytes (float32)} = 102,400,000 \text{ bytes}$
- In MB:  $102,400,000 / (1024 \times 1024) = 97.7 \text{ MB}$

This is why vocabulary size dominates the deployment budget for small-model, large-vocab systems.

### Q2: Positional Encoding Memory

Compare memory requirements for learned vs sinusoidal positional encoding with `max_seq_len=2048` and `embed_dim=512`.

#### 💡 Answer

**Learned PE:**  $2,048 \times 512 \times 4 = 4,194,304 \text{ bytes} = 4.0 \text{ MB}$  (1,048,576 parameters)

**Sinusoidal PE:** 0 bytes (0 parameters — computed from a closed-form formula)

For large models, learned PE adds real memory. GPT-3's learned PE costs an additional 96 MB ( $2048 \times 12288 \times 4 \text{ bytes}$ ). Models that need cheaper or longer-context PE often switch to sinusoidal or rotary variants.

### Q3: Lookup Complexity

What is the time complexity of looking up embeddings for a batch of 32 sequences, each with 128 tokens?

#### 💡 Answer

**O(1) per token**, or **O(batch\_size × seq\_len)** =  $O(32 \times 128) = O(4,096)$  total.

Lookup is constant time per token because it is direct array indexing: `weight[token_id]`. For 4,096 tokens you do 4,096 constant-time loads.

Vocabulary size does **not** affect lookup time. Looking up from a 1,000-word table is the same speed as from a 100,000-word table (modulo cache effects). It is indexing, not search.

### Q4: Embedding Dimension Scaling

You have an embedding layer with `vocab_size=50000`, `embed_dim=512` using 98 MB. If you double `embed_dim` to 1024, what happens to memory?

 **Answer**

Memory **doubles to 195 MB**.

Embedding memory scales linearly with embedding dimension:

- Original:  $50,000 \times 512 \times 4 = 98 \text{ MB}$
- Doubled:  $50,000 \times 1,024 \times 4 = 195 \text{ MB}$

Each doubling of `embed_dim` doubles both storage and memory bandwidth — the reason production models balance dimension against available memory rather than maxing it out.

**Q5: Sinusoidal Extrapolation**

You trained a model with sinusoidal positional encoding and `max_seq_len=512`. Can you process sequences of length 1024 at inference time? What about with learned positional encoding?

 **Answer**

**Sinusoidal PE: Yes** - can extrapolate to length 1024 (or any length)

The mathematical formula creates unique encodings for any position. You simply compute:

```
pe_1024 = create_sinusoial_embeddings(max_seq_len=1024, embed_dim=512)
```

**Learned PE: No** - cannot handle sequences longer than training maximum

Learned PE creates a fixed embedding table of shape `(max_seq_len, embed_dim)`. For positions beyond 512, there are no learned embeddings. You must either: - Retrain with larger `max_seq_len` - Interpolate position embeddings (advanced technique) - Truncate sequences to 512 tokens

This is why many production models use sinusoidal or relative positional encodings that can handle variable lengths.

## 19.11 Key Takeaways

- **Embedding lookup is sparse, not dense:** rows are indexed directly ( $O(1)$  per token), only accessed rows receive gradients, and no one-hot tensor is ever materialized.
- **Positional encoding is additive geometry, not metadata:** learned PE is flexible but length-capped; sinusoidal PE extrapolates but is fixed — same shape, very different deployment behavior.
- **Memory scales linearly with `vocab`  $\times$  `dim`:** doubling either component doubles both storage and memory-bandwidth traffic, which is why dimension choices are always weighed against HBM budget.
- **Embeddings are cheap to compute and expensive to store:** they take ~10–15% of training time but 30–40% of model memory in large-vocab models — the opposite profile of attention.

**Coming next:** Module 12 builds scaled dot-product attention, the mechanism that turns the static `(B, S, D)` embedding tensor from this module into a context-aware representation where every position has been shaped by every other.

## 19.12 Further Reading

While embedding tables may seem like simple lookup dictionaries, their architectural design profoundly impacts system-level performance, memory bandwidth constraints, and distributed training topologies. For students seeking to understand the mathematical underpinnings and the hardware-software co-design of these continuous representations, the following milestones are critical:

### 19.12.1 Seminal Papers

- **Word2Vec** - Mikolov et al. (2013). Introduced efficient learned word embeddings through context prediction. Though your implementation learns embeddings end-to-end, Word2Vec established the paradigm that semantic similarity should be encoded spatially. **Systems Implication:** By replacing massive, infinitely sparse one-hot encoded vectors with dense, low-dimensional continuous arrays, this shift dramatically compressed the memory footprint. It transformed textual analysis from cache-thrashing sparse lookups into dense matrix multiplications, fundamentally pushing the operations toward the compute-bound ceiling of the Roofline model. [arXiv:1301.3781](#)
- **Attention Is All You Need** - Vaswani et al. (2017). Introduced sinusoidal positional encoding and demonstrated that learned embeddings combined with positional information enable powerful, highly parallel sequence models. **Systems Implication:** Replaced recursive token-by-token processing with self-attention over continuous embeddings, entirely breaking the sequential compute bottleneck of RNNs and allowing massive, unhindered data-parallelization across modern GPU streaming multi-processors. [arXiv:1706.03762](#)
- **BERT: Pre-training of Deep Bidirectional Transformers** - Devlin et al. (2018). Demonstrated how token embeddings perfectly fuse with positional and segment encodings for deep language understanding. **Systems Implication:** Proved the empirical value of scaling parameter counts to hundreds of millions, thereby exposing the limits of single-chip High Bandwidth Memory (HBM) and necessitating multi-GPU tensor parallel paradigms just to host the vast embedding tables and model states. [arXiv:1810.04805](#)

### 19.12.2 Additional Resources

- **Blog Post:** “[The Illustrated Word2Vec](#)” by Jay Alammar - Visual explanation of learned word embeddings and semantic relationships
- **Documentation:** [PyTorch nn.Embedding](#) - See production embedding implementation
- **Paper:** “[GloVe: Global Vectors for Word Representation](#)” - Pennington et al. (2014) - Alternative embedding approach based on co-occurrence statistics

## 19.13 What's Next

You now have static geometry: every token ID maps to a fixed vector in space, with positional information layered on top. But the embedding for “bank” is the same whether the next word is “river” or “loan” — the representation has no context. That is the question Module 12 answers.

### **i** Coming Up: Module 12 — Attention

You'll build scaled dot-product attention: the mechanism that lets every embedding look at every other embedding in the sequence and re-weight itself accordingly. Your static  $(B, S, D)$  embedding tensor becomes a *context-aware*  $(B, S, D)$  tensor where each position has been shaped by the rest of the sentence.

### Where these embeddings show up downstream:

Table 19.7 traces how this module is reused by later parts of the curriculum.

Table 19.7: How embeddings feed into attention and transformer modules.

| Module                  | What It Does                    | Your Embeddings In Action                                              |
|-------------------------|---------------------------------|------------------------------------------------------------------------|
| <b>12: Attention</b>    | Context-aware representations   | <code>attention(embed_layer(tokens))</code> produces query, key, value |
| <b>13: Transformers</b> | Full sequence-to-sequence stack | <code>transformer(embed_layer(src), embed_layer(tgt))</code>           |

## 19.14 Get Started

### 💡 Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### ⚠️ Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.

—————| **12: Attention** | Context-aware representations | `attention(embed_layer(tokens))` creates query, key, value || **13: Transformers** | Complete sequence-to-sequence | `transformer(embed_layer(src), embed_layer(tgt))` |



## 🔥 Chapter 20

# Module 12: Attention

Attention's  $O(N^2)$  compute cost is only half the story. Its  $O(N^2)$  memory footprint for the scores matrix and  $O(N)$  KV-cache growth during decoding are what actually limit real inference workloads, turning attention from a FLOP problem into an HBM bandwidth problem. This module builds scaled dot-product attention, multi-head attention, and causal masking from scratch so you can see exactly where the quadratic wall comes from before FlashAttention and friends try to climb it.

**i** Module Info

**ARCHITECTURE TIER** | Difficulty: ●●●○ | Time: 5-7 hours | Prerequisites: 01-08, 10-11

**Prerequisites:** Modules 01-08 and 10-11 means you should understand:

- Tensor operations and shape manipulation (Module 01)
- Activations, particularly softmax (Module 02)
- Linear layers and weight projections (Module 03)
- Autograd for gradient computation (Module 06)
- Tokenization and embeddings (Modules 10-11)

If you can explain why `softmax(x).sum(axis=-1)` equals 1.0 and how embeddings convert token IDs to dense vectors, you're ready.

## 20.1 Overview

Attention is the mechanism behind GPT, BERT, and every modern LLM. In this module you build it from scratch — scaled dot-product attention and multi-head attention — the same math that runs in production transformers, written in NumPy so you can read every line.

The shift attention introduced is simple to state. RNNs squeeze a whole sequence into one fixed-size hidden state and hope nothing important falls out. Attention lets every position read directly from every other position, weighted by relevance computed on the fly. That's it. The cost of that freedom is the equation you'll implement:  $\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k}) V$ . The  $QK^T$  term creates an  $n \times n$  matrix — quadratic in sequence length. By the end of this module you'll have written that matrix, watched it dominate memory at long context, and understood exactly why FlashAttention and friends exist.

## 20.2 Learning Objectives

💡 By completing this module, you will:

- **Implement** scaled dot-product attention with vectorized operations that reveal  $O(n^2)$  memory complexity
- **Build** multi-head attention for parallel processing of different relationship types across representation subspaces

- **Master** attention weight computation, normalization, and the query-key-value paradigm
- **Understand** quadratic memory scaling and why attention becomes the bottleneck in long-context transformers
- **Connect** your implementation to production frameworks and understand why efficient attention research matters at scale

## 20.3 What You'll Build




Figure 20.1: **Attention Mechanism:** Information retrieval paradigm where similarity scores between Queries and Keys determine how much of each Value is retrieved.

### Implementation roadmap:

Table 20.1 lays out the implementation in order, one part at a time.

Table 20.1: **Implementation roadmap for scaled dot-product and multi-head attention.**

| Part | What You'll Implement                       | Key Concept                                         |
|------|---------------------------------------------|-----------------------------------------------------|
| 1    | <code>scaled_dot_product_attention()</code> | Core attention mechanism with $QK^T$ similarity     |
| 2    | Attention weight normalization              | Softmax converts scores to probability distribution |
| 3    | Causal masking support                      | Preventing attention to future positions            |
| 4    | <code>MultiHeadAttention.__init__()</code>  | Linear projections and head configuration           |
| 5    | <code>MultiHeadAttention.forward()</code>   | Split, attend, concatenate pattern                  |

### The pattern you'll enable:

```

# Multi-head attention for sequence processing
mha = MultiHeadAttention(embed_dim=512, num_heads=8)
output = mha(embeddings, mask) # Learn different relationship types in parallel

```

### 20.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Full transformer blocks (that's Module 13: Transformers)
- Positional encoding (you built this in Module 11: Embeddings)
- Efficient attention variants like FlashAttention (production optimization beyond scope)
- Cross-attention for encoder-decoder models (PyTorch does this with separate Q vs K/V inputs)

**You are building the core attention mechanism.** Complete transformer architectures come next.

## 20.4 API Reference

This section provides a quick reference for the attention functions and classes you'll build. Use this as your implementation guide and debugging reference.

### 20.4.1 Scaled Dot-Product Attention Function

```
scaled_dot_product_attention(Q, K, V, mask=None) -> (output, attention_weights)
```

Computes the fundamental attention operation that powers all transformers.

**Parameters:**

- **Q:** Query tensor (batch\_size, seq\_len, d\_model) — what each position is looking for
- **K:** Key tensor (batch\_size, seq\_len, d\_model) — what's available at each position
- **V:** Value tensor (batch\_size, seq\_len, d\_model) — actual content to retrieve
- **mask:** Optional (batch\_size, seq\_len, seq\_len) — 1.0 for allowed positions, 0.0 for masked

**Returns:**

- **output:** Attended values (batch\_size, seq\_len, d\_model)
- **attention\_weights:** Attention matrix (batch\_size, seq\_len, seq\_len) showing focus patterns

### 20.4.2 MultiHeadAttention Class

Multi-head attention runs multiple attention mechanisms in parallel, each learning to focus on different types of relationships.

**Constructor:**

```
MultiHeadAttention(embed_dim, num_heads) -> MultiHeadAttention
```

Creates multi-head attention with `embed_dim` // `num_heads` dimensions per head.

**Core Methods:**

Table 20.2 lists the methods on `MultiHeadAttention`.

Table 20.2: Core methods on the `MultiHeadAttention` class.

| Method                  | Signature                                       | Description                         |
|-------------------------|-------------------------------------------------|-------------------------------------|
| <code>forward</code>    | <code>forward(x, mask=None) -&gt; Tensor</code> | Apply multi-head attention to input |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code>    | Return all trainable parameters     |

**Attributes:**

- `embed_dim`: Total embedding dimension
- `num_heads`: Number of parallel attention heads
- `head_dim`: Dimension per head (`embed_dim // num_heads`)
- `q_proj, k_proj, v_proj`: Linear projections for queries, keys, values
- `out_proj`: Output linear layer to mix information across heads

## 20.5 Core Concepts

This section covers the fundamental ideas you need to understand attention deeply. These concepts apply to every transformer-based model in production today.

### 20.5.1 Query, Key, Value: The Information Retrieval Paradigm

Attention treats sequence processing as a soft database lookup. Three roles, all learned:

- **Query** — what this position is looking for.
- **Key** — what each other position advertises about itself.
- **Value** — what each position actually hands over if attended to.

To translate this into code, your implementation creates these three components using dedicated linear projections, mapping the input sequence into distinct representations for querying, indexing, and retrieving: The code in `?@lst-12-qkv-projections` makes this concrete.

```
# From MultiHeadAttention.__init__
self.q_proj = Linear(embed_dim, embed_dim) # Learn what to search for
self.k_proj = Linear(embed_dim, embed_dim) # Learn what to index by
self.v_proj = Linear(embed_dim, embed_dim) # Learn what to retrieve

# From MultiHeadAttention.forward
Q = self.q_proj.forward(x) # Transform input to queries
K = self.k_proj.forward(x) # Transform input to keys
V = self.v_proj.forward(x) # Transform input to values
```

: **Listing 12.1 — Query, key, value projections.** Three learned Linear layers turn one input into three role-specific tensors for attention. `{#lst-12-qkv-projections}`

Three projections, three roles, one input. The same embedding plays all three parts at every position — and what those parts mean is whatever the optimizer settles on.

### 20.5.2 Scaled Dot-Product Attention: Similarity as Relevance

The core computation asks one question: how similar is each query to each key? A dot product is the cheapest reasonable answer — large positive values mean the vectors point the same way, near-zero means orthogonal, large negative means opposite. For positions `i` and `j`, the score  $Q[i] \cdot K[j]$  is exactly that.

The  $1/\sqrt{d_k}$  factor exists for a specific, fixable failure. If the components of `Q` and `K` are roughly unit-variance, then  $Q[i] \cdot K[j]$  is a sum of  $d_k$  independent products — its variance grows linearly with  $d_k$ , so its standard deviation grows like  $\sqrt{d_k}$ . At  $d_k = 64$  the typical score is around 8; at  $d_k = 512$  it's around 23. Feed that to softmax and one entry takes essentially all the probability mass, while every other entry sits in the flat tail of  $\exp$  where the gradient is nearly zero. Training stalls. Dividing by  $\sqrt{d_k}$  cancels exactly that growth, holding score variance at roughly 1 regardless of head dimension.

Your implementation computes this using vectorized matrix operations:

The code in `?@lst-12-scaled-scores` makes this concrete.

```

# From scaled_dot_product_attention (lines 303-319)
d_model = Q.shape[-1]

# Compute all query-key similarities at once using matmul
# This is mathematically equivalent to nested loops computing Q[i] · K[j]
# for all i,j pairs, but vectorized for efficiency
K_t = K.transpose(-2, -1) # Transpose to align dimensions
scores = Q.matmul(K_t)    # (batch, seq_len, seq_len) - the O(n²) matrix

# Scale by 1/√d_k for numerical stability
scale_factor = 1.0 / math.sqrt(d_model)
scores = scores * scale_factor

```

: **Listing 12.2 — Scaled dot-product score matrix.** Computes  $QK^T / \sqrt{d_k}$ , producing the  $O(N^2)$  scores tensor that dominates attention memory. {#lst-12-scaled-scores}

The resulting `scores` tensor is the attention matrix before normalization. Element `[i, j]` represents how much position `i` should attend to position `j`. The vectorized `matmul` operation computes all  $n^2$  query-key pairs simultaneously—while much faster than Python loops, it still creates the full  $O(n^2)$  attention matrix that dominates memory usage at scale.

### 20.5.3 Attention Weights and Softmax Normalization

Raw similarity scores need to become a probability distribution. Softmax transforms scores into positive values that sum to 1.0 along each row, creating a proper weighted average. This ensures that for each query position, the attention weights over all key positions form valid mixing coefficients.

The softmax operation  $\exp(\text{scores}[i, j]) / \sum_k \exp(\text{scores}[i, k])$  has important properties. It's differentiable, allowing gradients to flow during training. It amplifies differences: a score of 2.0 becomes much more prominent than 1.0 after exponentiation. And it's translation-invariant: adding the same constant to all scores doesn't change the output (exploited for numerical stability).

Here's the complete attention weight computation with masking support:

The code in `?@lst-12-masked-softmax` makes this concrete.

```

# From scaled_dot_product_attention

# Apply causal mask if provided (set masked positions to large negative)
if mask is not None:
    mask_data = mask.data
    adder_mask = (1.0 - mask_data) * MASK_VALUE # MASK_VALUE = -1e9
    adder_mask_tensor = Tensor(adder_mask, requires_grad=False)
    scores = scores + adder_mask_tensor

# Softmax converts scores to probability distribution
softmax = Softmax()
attention_weights = softmax(scores, dim=-1) # Normalize along last dimension

# Apply to values: weighted combination
output = attention_weights.matmul(V)

```

: **Listing 12.3 — Masked softmax and value-weighted sum.** Adds  $-1e9$  at masked positions, normalizes via softmax, then weights values to produce the attended output. {#lst-12-masked-softmax}

The mask trick is worth pausing on. Masked positions get  $-1e9$  added to their score; after `exp`, that's a number indistinguishable from zero, so they receive no attention weight at all. Allowed positions get  $+0$  and pass through unchanged. The whole operation stays differentiable — there is no `if` in the gradient path — and the constraint is enforced exactly.

### 20.5.4 Multi-Head Attention: Parallel Relationship Learning

A single attention head learns one similarity function — one notion of “what counts as relevant”. Real sequences have several at once: local syntactic agreement, long-range coreference, positional offsets, semantic similarity. Multi-head attention runs several attention mechanisms side by side, each with its own learned Q/K/V projections, so each head can specialize.

The crucial design choice is to *split* the embedding dimension across heads, not duplicate it. With `embed_dim=512` and `num_heads=8`, each head operates on  $512/8 = 64$  dimensions. Parameter count stays the same as a single 512-dim head, but you get 8 narrower attention mechanisms running in parallel instead of one wide one. In practice, heads specialize on their own — visualizations of trained models show one head tracking adjacent tokens, another tracking subject-verb pairs across long distances, others doing things no one has named.

Your implementation handles this through reshape and transpose operations:

The code in `?@lst-12-multihead-forward` makes this concrete.

```
# From MultiHeadAttention.forward

# Project to Q, K, V (each is batch, seq, embed_dim)
Q = self.q_proj.forward(x)
K = self.k_proj.forward(x)
V = self.v_proj.forward(x)

# Reshape to separate heads: (batch, seq, num_heads, head_dim)
Q = Q.reshape(batch_size, seq_len, self.num_heads, self.head_dim)
K = K.reshape(batch_size, seq_len, self.num_heads, self.head_dim)
V = V.reshape(batch_size, seq_len, self.num_heads, self.head_dim)

# Transpose to (batch, num_heads, seq, head_dim) for parallel processing
Q = Q.transpose(1, 2)
K = K.transpose(1, 2)
V = V.transpose(1, 2)

# Apply attention to all heads at once
attended, _ = scaled_dot_product_attention(Q, K, V, mask=mask_reshaped)

# Transpose back and concatenate heads
attended = attended.transpose(1, 2) # (batch, seq, num_heads, head_dim)
concat_output = attended.reshape(batch_size, seq_len, self.embed_dim)

# Mix information across heads with output projection
output = self.out_proj.forward(concat_output)
```

: **Listing 12.4 — Multi-head attention forward pass.** Reshape splits the embedding across heads, attention runs in parallel per head, then reshape and `out_proj` recombine. `{#lst-12-multihead-forward}`

The reshape-transpose-attend-transpose-reshape dance separates heads for independent processing, then recombines their outputs. The final output projection learns how to mix information discovered by different heads, creating a rich representation that captures multiple relationship types simultaneously.

### 20.5.5 Causal Masking: Preventing Information Leakage

GPT generates one token at a time, left to right. To train it, we feed in a whole sentence at once and ask it to predict the next token at *every* position in parallel — but only if each position is forbidden from peeking at the tokens that come after it. Otherwise the task is trivial: position 5 just copies token 6 and reports a perfect prediction. Causal masking is what makes the parallel training match the sequential generation: position  $i$  may attend only to positions  $0 \dots i$ .

The mask itself is a lower-triangular matrix — ones on and below the diagonal, zeros above:

```
[[1, 0, 0, 0], # Position 0 can only see itself
 [1, 1, 0, 0], # Position 1 sees 0 and 1
 [1, 1, 1, 0], # Position 2 sees 0, 1, 2
 [1, 1, 1, 1]] # Position 3 sees the whole prefix
```

Combined with the  $-1e9$  trick from the previous section, exactly the upper triangle of every attention matrix is zeroed out before the weighted sum. The reader should pause on this: the mask is the *only* thing that turns a bidirectional encoder (like BERT) into an autoregressive decoder (like GPT). Same code, same parameters, one extra triangular tensor.

### 20.5.6 Computational Complexity: The $O(n^2)$ Reality

Attention’s power comes from all-to-all connectivity: every position can attend to every other position. But this creates quadratic scaling in both computation and memory. For sequence length  $n$ , the attention matrix has  $n^2$  elements. The vectorized  $Q @ K^T$  operation computes all  $n^2$  similarity scores in one matrix multiplication, softmax normalizes  $n^2$  values, and applying attention to values multiplies  $n^2$  weights by the value vectors.

The memory cost is the part that bites first. For GPT-3 with 2048-token context, a single attention matrix stores  $2048^2 = 4,194,304$  float32 values — 16 MB. With 96 layers stacked, attention matrices alone consume 1.5 GB before you account for activations, gradients, Q/K/V projections, or anything else. That is the quadratic wall, and it is the reason every long-context system you’ve heard of has a paper attached to it.

Table 20.3 summarises the time and memory cost of the core operations.

Table 20.3: Time and memory complexity of the three attention computation steps.

| Operation    | Time Complexity                     | Memory Complexity          | Dominates When                  |
|--------------|-------------------------------------|----------------------------|---------------------------------|
| $QK^T$       | $O(n^2 \times d)$                   | $O(n^2)$                   | Long sequences                  |
| Softmax      | $O(n^2)$                            | $O(n^2)$                   | Always stores full matrix       |
| Weights @ V  | $O(n^2 \times d)$                   | $O(n \times d)$            | Output reuses attention weights |
| <b>Total</b> | <b><math>O(n^2 \times d)</math></b> | <b><math>O(n^2)</math></b> | $n > d$ (long sequences)        |

For comparison, feed-forward networks in transformers have  $O(n \times d^2)$  complexity. When sequence length  $n$  exceeds embedding dimension  $d$  (common in modern models), attention’s  $O(n^2)$  term dominates, making it the primary bottleneck. This explains why research into efficient attention variants like sparse attention, linear attention, and FlashAttention is crucial for production systems.

**i Systems Implication: The Memory Wall (FlashAttention)**

The  $O(n^2)$  memory requirement of attention is the single biggest bottleneck in modern generative AI. It dictates the “context window” limit of every LLM. To calculate attention, standard implementations must write the massive  $N \times N$  similarity matrix to the GPU’s slow High Bandwidth Memory (HBM) and then read it back for the softmax operation. Systems engineers bypass this wall using **FlashAttention**. This algorithm doesn’t just compute attention in small “blocks” that fit entirely inside the GPU’s ultra-fast, but tiny, SRAM; crucially, it **fuses** the operations to avoid the costly **HBM read/write round-trips** entirely.

## 20.6 Common Errors

These are the errors you’ll encounter most often when implementing attention. Understanding them will save hours of debugging.

### 20.6.1 Shape Mismatch in Attention

**Error:** `ValueError: Cannot perform matrix multiplication: (2, 4, 64) @ (2, 4, 64). Inner dimensions must match`

When computing  $Q @ K^T$ , the key tensor needs transposing. The matrix multiplication  $Q @ K$  has shape  $(batch, seq\_len, d\_model) @ (batch, seq\_len, d\_model)$ , which fails because the inner dimensions are both  $d\_model$ . You need  $Q @ K.transpose()$  to get  $(batch, seq\_len, d\_model) @ (batch, d\_model, seq\_len)$ , producing the correct  $(batch, seq\_len, seq\_len)$  attention matrix.

**Fix:** Always transpose K before the matmul: `scores = Q.matmul(K.transpose(-2, -1))`

### 20.6.2 Attention Weights Don’t Sum to 1

**Error:** `AssertionError: Attention weights don't sum to 1`

This happens when softmax is applied to the wrong axis. Attention weights must form a probability distribution over key positions for each query position. If you apply softmax along the wrong dimension, you’ll get values that don’t sum to 1.0 per row.

**Fix:** Use `softmax(scores, dim=-1)` to normalize along the last dimension (across keys for each query)

### 20.6.3 Multi-Head Dimension Mismatch

**Error:** `ValueError: embed_dim (512) must be divisible by num_heads (7)`

Multi-head attention splits the embedding dimension across heads. If `embed_dim=512` and `num_heads=7`, you’d get  $512/7=73.14$  dimensions per head, which doesn’t work with integer tensor shapes. The architecture requires exact divisibility.

**Fix:** Choose `num_heads` that evenly divides `embed_dim`. Common pairs:  $(512, 8)$ ,  $(768, 12)$ ,  $(1024, 16)$

### 20.6.4 Mask Broadcasting Errors

**Error:** `ValueError: operands could not be broadcast together with shapes (2,1,4,4) (2,4,4)`

Multi-head attention expects masks with a head dimension. If you pass a 3D mask  $(batch, seq, seq)$  but the implementation expects 4D  $(batch, heads, seq, seq)$ , broadcasting fails. The mask needs reshaping to add a dimension that broadcasts across all heads.

**Fix:** Reshape mask: `mask.reshape(batch, 1, seq_len, seq_len)` to broadcast over heads

### 20.6.5 Gradient Flow Issues

**Error:** Loss doesn’t decrease during training despite correct forward pass



This can happen if you create new Tensor objects incorrectly, breaking the autograd graph. When applying masks or performing intermediate computations, ensure tensors maintain `requires_grad` appropriately.

**Fix:** Check that operations preserve gradient flow: `Tensor(result, requires_grad=True)` when needed

## 20.7 Production Context

### 20.7.1 Your Implementation vs. PyTorch

Your TinyTorch attention and PyTorch's `nn.MultiheadAttention` implement the same mathematical operations. The differences are in implementation efficiency, features, and flexibility. PyTorch uses highly optimized C++ kernels, supports additional attention variants, and integrates with production training systems.

Table 20.4 places your implementation side by side with the production reference for direct comparison.

Table 20.4: Feature comparison between TinyTorch attention and `nn.MultiheadAttention`.

| Feature                    | Your Implementation              | PyTorch                                        |
|----------------------------|----------------------------------|------------------------------------------------|
| <b>Core Algorithm</b>      | Scaled dot-product attention     | Same mathematical operation                    |
| <b>Multi-Head</b>          | Split-attend-concat pattern      | Identical architecture                         |
| <b>Backend</b>             | NumPy (Python loops)             | C++ CUDA kernels                               |
| <b>Speed</b>               | 1x (baseline)                    | 50-100x faster on GPU                          |
| <b>Memory Optimization</b> | Stores full attention matrix     | Optional FlashAttention integration            |
| <b>Batch First</b>         | <code>(batch, seq, embed)</code> | Configurable via <code>batch_first=True</code> |
| <b>Cross-Attention</b>     | Self-attention only              | Separate Q vs K/V inputs supported             |
| <b>Key Padding Mask</b>    | Manual mask creation             | Built-in mask utilities                        |

### 20.7.2 Code Comparison

The following comparison shows equivalent attention operations in TinyTorch and PyTorch. Notice how the high-level API and shape conventions match almost exactly.

## 20.8 Your TinyTorch

```
from tinytorch.core.attention import MultiHeadAttention
from tinytorch.core.tensor import Tensor
import numpy as np

# Create multi-head attention
mha = MultiHeadAttention(embed_dim=512, num_heads=8)
```

```

# Input embeddings (batch=2, seq=10, dim=512)
x = Tensor(np.random.randn(2, 10, 512))

# Apply attention
output = mha.forward(x) # (2, 10, 512)

# With causal masking
mask = Tensor(np.tril(np.ones((2, 10, 10))))
output_masked = mha.forward(x, mask)

```

## 20.9 PyTorch

```

import torch
import torch.nn as nn

# Create multi-head attention
mha = nn.MultiheadAttention(embed_dim=512, num_heads=8,
                           batch_first=True)

# Input embeddings (batch=2, seq=10, dim=512)
x = torch.randn(2, 10, 512)

# Apply attention (PyTorch returns output + weights)
output, weights = mha(x, x, x) # Self-attention: Q=K=V=x

# With causal masking (upper triangle = -inf)
mask = torch.triu(torch.ones(10, 10) * float('-inf'), diagonal=1)
output_masked, _ = mha(x, x, x, attn_mask=mask)

```

Let's walk through the key differences:

- **Line 1-2 (Imports):** TinyTorch separates attention into its own module; PyTorch includes it in `torch.nn`. Both follow modular design patterns.
- **Line 4-5 (Construction):** API is nearly identical. PyTorch adds `batch_first=True` for compatibility with older code that expected (seq, batch, embed) order.
- **Line 8 (Input):** Shape conventions match exactly: (batch, seq, embed). This is the modern standard.
- **Line 11 (Forward Pass):** TinyTorch uses `mha.forward(x)` with `x` as both Q, K, V (self-attention). PyTorch makes this explicit with `mha(x, x, x)`, allowing cross-attention where Q differs from K/V.
- **Line 14-15 (Masking):** TinyTorch uses 0/1 masks (0=masked). PyTorch uses additive masks (-inf=masked). Both work, but PyTorch's convention integrates better with certain optimizations.

### 💡 What's Identical

The mathematical operations, architectural patterns, and shape conventions are identical. Multi-head attention works the same way in production. Understanding your implementation means understanding PyTorch's attention.

### 20.9.1 Why Attention Matters at Scale

To appreciate why attention research is crucial, consider the scaling characteristics of modern language models:

- **GPT-3** (96 layers, 2048 context): ~1.5 GB just for attention matrices during the forward pass, ~7.5 GB once gradients and optimizer state are added during training
- **GPT-4** (estimated 120 layers, 32K context): would require ~480 GB for attention alone without optimization, far exceeding any single-GPU memory budget
- **Long-context models** (100K+ tokens): attention becomes computationally prohibitive without algorithmic improvements

These constraints drive modern attention research:

- **FlashAttention**: Reformulates computation to reduce memory from  $O(n^2)$  to  $O(n)$  without approximation, enabling 8x longer contexts
- **Sparse Attention**: Only compute attention for specific patterns (local windows, strided access), reducing complexity to  $O(n \log n)$  or  $O(n\sqrt{n})$
- **Linear Attention**: Approximate attention with linear complexity  $O(n)$ , trading accuracy for scale
- **State Space Models**: Alternative architectures (Mamba, RWKV) that avoid attention's quadratic cost entirely

The attention mechanism you built is mathematically identical to production systems, but the  $O(n^2)$  wall explains why so much research focuses on making it tractable at scale.

#### 💡 Check Your Understanding — Attention

Before moving on, verify you can articulate each of the following:

- Why the  $O(N^2)$  attention bottleneck is fundamentally **memory-bound**, not compute-bound — the full  $N \times N$  scores matrix has to be written to and read from HBM — and exactly what FlashAttention does to fuse softmax with the matmuls in SRAM to sidestep it.
- How the  $1/\sqrt{d_k}$  scaling factor keeps softmax out of its saturated tail, and why skipping it causes training to stall (gradients collapse to zero for all but one key per query).
- Why splitting `embed_dim` across heads ( $8 \times 64$  instead of  $1 \times 512$ ) preserves parameter count but delivers representation diversity, cache locality, and parallelism — three system-level wins from one architectural choice.
- How causal masking turns a bidirectional encoder into an autoregressive decoder with a single triangular tensor — and why, in a naive implementation, it saves zero computation even though it zeroes out half the matrix.

If any of these feels fuzzy, revisit the *Core Concepts* section (especially *Computational Complexity: The  $O(n^2)$  Reality*) before moving on.

## 20.10 Self-Check Questions

Test yourself with these systems thinking questions. They're designed to build intuition for the performance characteristics you'll encounter in production ML.

### Q1: Memory Calculation

For sequence length 1024, how much memory does a single attention matrix require (float32)? What about sequence length 2048?

 **Answer****Sequence length 1024:**

- Attention matrix:  $1024 \times 1024 = 1,048,576$  elements
- Memory:  $1,048,576 \times 4 \text{ bytes} = \mathbf{4.0 \text{ MB}}$

**Sequence length 2048:**

- Attention matrix:  $2048 \times 2048 = 4,194,304$  elements
- Memory:  $4,194,304 \times 4 \text{ bytes} = \mathbf{16.0 \text{ MB}}$

**Scaling factor:** Doubling sequence length quadruples memory ( $2^2 = 4\times$ )

For GPT-3 (96 layers, 2048 context):

- $96 \text{ layers} \times 16.0 \text{ MB} = \mathbf{1.5 \text{ GB}}$  just for attention matrices.
- This excludes Q/K/V projections, gradients, and every other tensor.

**Q2: Attention Bottleneck**

A transformer layer has attention ( $O(n^2 \times d)$ ) and feed-forward network ( $O(n \times d^2)$ ). For `embed_dim=512`, at what sequence length does attention dominate?

 **Answer****Complexity comparison:**

- Attention:  $O(n^2 \times d) = O(n^2 \times 512)$
- FFN:  $O(n \times d^2) = O(n \times 512^2) = O(n \times 262,144)$

**Crossover point:**  $n^2 \times 512 > n \times 262,144$

- Simplify:  $n > 262,144 / 512 = \mathbf{512}$

**When  $n > 512$ ,** attention becomes the memory bottleneck.

**Real-world implications:**

- Short sequences ( $n=128$ ): FFN dominates, 262K vs 8K operations
- Medium sequences ( $n=512$ ): break-even point
- Long sequences ( $n=2048$ ): attention dominates, 2M vs 262K operations
- **This is why GPT-3 (2048 context) needed attention optimization.**

**Q3: Multi-Head Efficiency**

Why use 8 heads of 64 dimensions instead of 1 head of 512 dimensions? Parameters are the same—what's the systems difference?

 **Answer****Parameter count (both are identical):**

- 8 heads  $\times$  64 dims: `Linear(512→512)` for Q, K, V, Out =  $4 \times (512 \times 512 + 512)$  weights+biases
- 1 head  $\times$  512 dims: same projection parameters

**Key differences:****1. Parallelization:**

- 8 heads can process in parallel on modern GPUs (separate CUDA streams)
- Each head's smaller matmul operations utilize GPU cores more efficiently

**2. Representation diversity:**

- 8 heads learn 8 different similarity functions (syntax, semantics, position, etc.)

- 1 head learns a single monolithic similarity function
- Training discovers specialization automatically

### 3. Cache efficiency:

- Smaller head\_dim (64) fits better in GPU cache and shared memory
- A single 512-dim head causes more cache misses

### 4. Gradient flow:

- Multiple heads provide diverse gradient signals during backpropagation
- A single head has one gradient path, slower learning

**Empirical result:** 8 heads consistently outperform 1 head at equal parameter count. Diversity is the whole point.

## Q4: Causal Masking Computation

Causal masking zeros out the upper triangle (roughly half the attention matrix). Do we save computation, or just ensure correctness?

### 💡 Answer

#### In your implementation: NO computation saved

Your code computes the full attention matrix, then adds  $-1e9$  to masked positions:

```
scores = Q.matmul(K_t) # Full  $n^2$  computation
scores = scores + adder_mask_tensor # Masking happens after
```

#### Why no savings:

- `Q.matmul(K_t)` computes all  $n^2$  scores
- Masking only affects softmax, not the initial computation
- We still store and normalize the full matrix

#### To actually save computation, you'd need:

1. Sparse matrix multiplication (skip masked positions in matmul)
2. Computing only the lower triangle of scores
3. Specialized CUDA kernels that exploit sparsity

#### Production optimizations:

- PyTorch's standard attention also computes the full matrix (same as yours)
- FlashAttention uses tiling to avoid materializing the full matrix but doesn't exploit sparsity
- Sparse attention (BigBird, Longformer) actually skips computation for sparse patterns

**Memory could be saved** by storing only the lower triangle ( $n^2/2$  elements), but it requires custom indexing.

## Q5: Gradient Memory

Training attention requires storing activations for backpropagation. How much memory does training need compared to inference?

### 💡 Answer

#### Forward pass (inference):

- Attention matrix:  $n^2$  values

**Backward pass (training) additional memory:**

- Gradient of attention weights:  $n^2$  values
- Gradient of Q, K, V:  $3 \times (n \times d)$  values
- Intermediate gradients from softmax:  $n^2$  values

**With Adam optimizer (standard for transformers):**

- First moment (momentum):  $n^2$  values
- Second moment (velocity):  $n^2$  values

**Total multiplier for attention matrix alone:**

- Forward:  $1 \times$  (attention weights)
- Backward:  $+2 \times$  (gradients)
- Optimizer:  $+2 \times$  (Adam state)
- **Total:  $5 \times$  inference memory**

**For GPT-3 scale (96 layers, 2048 context):**

- Inference:  $96 \times 16 \text{ MB} = 1.5 \text{ GB}$
- Training:  $96 \times 16 \text{ MB} \times 5 = 7.5 \text{ GB}$  just for attention gradients and optimizer state.

This excludes Q/K/V matrices, feed-forward networks, embeddings, and activations from other layers. Full GPT-3 training requires 350+ GB.

## 20.11 Key Takeaways

- **Attention is Q/K/V as a soft database lookup:** three learned projections of the *same* input let every position query every other, with similarity scores becoming mixing weights via softmax.
- **The  $O(N^2)$  cost is a memory wall, not a FLOP wall:** the  $QK^T$  matrix must be written to HBM before softmax reads it back — a round-trip that FlashAttention eliminates with SRAM-resident tiling.
- **Multi-head attention is diversity via splitting:** partitioning `embed_dim` across heads keeps parameters constant but yields cache-friendly, parallelizable heads that specialize during training.
- **Causal masking is what distinguishes GPT from BERT:** one triangular tensor turns the same parameters and same code into an autoregressive decoder by forbidding every position from peeking at future tokens.

**Coming next:** Module 13 wraps attention in the scaffolding that makes it trainable at depth — layer normalization, pre-norm residual streams, and  $4 \times$  MLPs — stacked into a full GPT.

## 20.12 Further Reading

The theoretical brilliance of attention is matched only by the engineering challenges required to deploy it at scale. For students who want to understand the architectural turning points and hardware-software co-design that conquered the  $O(n^2)$  memory wall:

### 20.12.1 Seminal Papers

- **Attention Is All You Need** - Vaswani et al. (2017). The paper that introduced transformers and the multi-head attention mechanism you just built. Shows how attention alone, without recurrence, achieves state-of-the-art results. **Systems Implication:** By completely discarding recurrence, it shattered the sequential compute bottleneck of RNNs, allowing entire sequences to be ingested in parallel, perfectly saturating the massive SIMD grids of modern GPUs. [arXiv:1706.03762](https://arxiv.org/abs/1706.03762)

- **BERT: Pre-training of Deep Bidirectional Transformers** - Devlin et al. (2018). Demonstrates how bidirectional attention (no causal mask) enables powerful language understanding. **Systems Implication:** Deep bidirectional attention required retaining massive  $N \times N$  activation checkpoints during the forward pass for backpropagation. This aggressively pushed the boundaries of GPU HBM, driving the industry-wide adoption of activation recomputation (gradient checkpointing) to trade compute for memory. [arXiv:1810.04805](#)
- **Language Models are Unsupervised Multitask Learners (GPT-2)** - Radford et al. (2019). Shows how causal attention with the masking pattern you implemented enables autoregressive language modeling at unprecedented scale. **Systems Implication:** As model footprints exploded beyond the capacity of a single GPU, autoregressive generation became heavily latency-bound. This architecture popularized large-scale tensor parallelism, splitting the attention heads across distinct hardware accelerators to preserve inference speed. [OpenAI](#)
- **FlashAttention: Fast and Memory-Efficient Exact Attention** - Dao et al. (2022). Directly addresses the  $O(n^2)$  memory bottleneck you experienced, achieving 2-4 $\times$  speedups without mathematical approximation. **Systems Implication:** By leveraging precise hardware-aware SRAM tiling, it entirely eliminated the need to materialize the massive attention matrix in slow HBM. This masterclass in systems engineering single-handedly rescued self-attention from the memory bandwidth ceiling, dragging it back to a compute-bound operation. [arXiv:2205.14135](#)

### 20.12.2 Additional Resources

- **Blog post:** “[The Illustrated Transformer](#)” by Jay Alammar - Visual explanations of attention mechanics that complement your implementation
- **Interactive tool:** [BertViz](#) - Visualize attention patterns in trained models to see the specialization you enabled with multi-head attention
- **Textbook:** “[Speech and Language Processing](#)” (Jurafsky & Martin, Chapter 9) - Formal treatment of attention in sequence-to-sequence models

## 20.13 What's Next

### Coming Up: Module 13 — Transformers

You have built the engine. Module 13 builds the chassis around it. The question Module 13 answers is: *what do you wrap attention in to actually train it?* Bare attention has two problems — its outputs collapse to similar values across positions (no per-token computation), and stacking it deeply makes gradients vanish. The transformer block fixes both, with feed-forward networks for per-position transformation, layer normalization to keep activations well-scaled, and residual connections to keep gradients flowing through arbitrarily many layers. Stack the result and you have GPT.

**Preview — how your attention gets used in future modules:**

Table 20.5 traces how this module is reused by later parts of the curriculum.

Table 20.5: **How attention feeds into the transformer block assembly.**

| Module                  | What It Does                | Your Attention In Action                                   |
|-------------------------|-----------------------------|------------------------------------------------------------|
| <b>13: Transformers</b> | Complete transformer blocks | <code>TransformerLayer(attention + FFN + LayerNorm)</code> |

| Module                  | What It Does         | Your Attention In Action                                                                            |
|-------------------------|----------------------|-----------------------------------------------------------------------------------------------------|
| <b>13: Transformers</b> | Residual connections | $x + \text{attention}(x)$ keeps gradients flowing                                                   |
| <b>13: Transformers</b> | Stacked layers       | $\text{attention} \rightarrow \text{FFN} \rightarrow \text{attention} \rightarrow \text{FFN} \dots$ |

## 20.14 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## 🔥 Chapter 21

# Module 13: Transformers

A transformer block's cost profile has two regimes that fight for the same HBM: attention at  $O(N^2)$  memory in sequence length, and MLPs at  $O(N \cdot d^2)$  compute in hidden width. Stack twelve blocks and the attention matrices alone can exceed the weight matrices once  $N$  crosses a few thousand tokens, which is why every production LLM ships with KV caching, activation checkpointing, and attention kernels tuned for the SRAM hierarchy. This module wires LayerNorm, MLPs, and causal self-attention into a working GPT so you can see exactly which components hit which wall first.

### Module Info

**ARCHITECTURE TIER** | Difficulty: ●●●● | Time: 8-10 hours | Prerequisites: 01-08, 10-12

You need tensors, layers, training loops, tokenization, embeddings, and attention already in place. If you can explain how multi-head attention turns queries, keys, and values into a weighted representation, you are ready for this chapter.

## 21.1 Overview

This is the chapter where everything snaps together. You have tensors, autograd, layers, a training loop, embeddings, and attention. In this module you wire them into a transformer block, stack the blocks, and end up with a working GPT — the same architecture that powers GPT, Claude, and LLaMA. By the end you can run `model.generate(prompt)` on something you wrote yourself.

A transformer block is a small recipe: layer-normalize, run multi-head attention, add a residual; layer-normalize, run an MLP, add a residual. That is it. Stack twelve of those between an embedding table and a language head, train on next-token prediction, and you have a 60M-parameter language model that produces coherent text.

The patterns you implement here — pre-norm, residual streams, causal masking,  $4 \times$  MLPs — are exactly what runs in production at billion-token scale. The optimizations differ; the architecture does not.

## 21.2 Learning Objectives

### By completing this module, you will:

- **Implement** layer normalization to stabilize training across deep networks with learnable scale and shift parameters
- **Design** complete transformer blocks combining self-attention, feed-forward networks, and residual connections using pre-norm architecture
- **Build** a full GPT model with token embeddings, positional encoding, stacked transformer blocks, and autoregressive generation

- **Analyze** parameter scaling and memory requirements, understanding why attention memory grows quadratically with sequence length
- **Master** causal masking to enable autoregressive generation while preventing information leakage from future tokens

## 21.3 What You'll Build

Figure 21.1 shows the full GPT stack you will assemble, from token IDs at the bottom to logits at the top.




Figure 21.1

### Implementation roadmap:

Table 21.1 lays out the implementation in order, one part at a time.

Table 21.1: **Implementation roadmap for LayerNorm, MLP, TransformerBlock, and GPT.**

| Step | What You'll Implement                                    | Key Concept                                          |
|------|----------------------------------------------------------|------------------------------------------------------|
| 1    | <code>LayerNorm</code> with learnable gamma/beta         | Stabilizes training by normalizing activations       |
| 2    | <code>MLP</code> with 4x expansion and GELU              | Provides non-linear transformation capacity          |
| 3    | <code>TransformerBlock</code> with pre-norm architecture | Combines attention and MLP with residual connections |

| Step | What You'll Implement                      | Key Concept                                  |
|------|--------------------------------------------|----------------------------------------------|
| 4    | GPT model with embeddings and blocks       | Complete autoregressive language model       |
| 5    | Autoregressive generation with temperature | Text generation with controllable randomness |

The pattern you'll enable:

```
# Building and using a complete language model
model = GPT(vocab_size=50000, embed_dim=768, num_layers=12, num_heads=12)
logits = model.forward(tokens) # Process input sequence
generated = model.generate(prompt, max_new_tokens=50) # Generate text
```

### 21.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- KV caching for efficient generation (production systems cache keys/values to avoid recomputation)
- FlashAttention or other memory-efficient attention (PyTorch uses specialized CUDA kernels)
- Mixture of Experts or sparse transformers (advanced scaling techniques)
- Multi-query or grouped-query attention (used in modern LLMs for efficiency)

**You are building the canonical transformer architecture.** Optimizations come later.

## 21.4 API Reference

This section documents the transformer components you'll implement. Each class builds on the previous, culminating in a complete language model.

### 21.4.1 Helper Functions

#### create\_causal\_mask

```
create_causal_mask(seq_len: int) -> Tensor
```

Creates a causal (autoregressive) attention mask that prevents positions from attending to future positions. Returns a lower triangular matrix where position  $i$  can only attend to positions  $j \leq i$ .

**Returns:** Tensor of shape  $(1, \text{seq\_len}, \text{seq\_len})$  with 1.0 for allowed positions, 0.0 for masked positions.

### 21.4.2 LayerNorm

```
LayerNorm(normalized_shape: int, eps: float = 1e-5) -> LayerNorm
```

Normalizes activations across features for each sample independently. Essential for stable training of deep transformer networks.

**Core Methods:**

Table 21.2 lists the methods on this class.

Table 21.2: Core methods on the LayerNorm class.

| Method                  | Signature                                    | Description                                                |
|-------------------------|----------------------------------------------|------------------------------------------------------------|
| <code>forward</code>    | <code>forward(x: Tensor) -&gt; Tensor</code> | Normalize across last dimension with learnable scale/shift |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code> | Returns <code>[gamma, beta]</code> learnable parameters    |

### 21.4.3 MLP (Multi-Layer Perceptron)

```
MLP(embed_dim: int, hidden_dim: int = None, dropout_prob: float = 0.1) -> MLP
```

Feed-forward network with 4x expansion, GELU activation, and projection back to original dimension.

**Core Methods:**

Table 21.3 lists the methods on this class.

Table 21.3: Core methods on the MLP feed-forward block.

| Method                  | Signature                                    | Description                                 |
|-------------------------|----------------------------------------------|---------------------------------------------|
| <code>forward</code>    | <code>forward(x: Tensor) -&gt; Tensor</code> | Apply Linear → GELU → Linear transformation |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code> | Returns weights and biases from both layers |

### 21.4.4 TransformerBlock

```
TransformerBlock(embed_dim: int, num_heads: int, mlp_ratio: int = 4, ff_dim: int = None, dropout_prob: float = 0.1) -> TransformerBlock
```

Complete transformer block with self-attention, MLP, layer normalization, and residual connections using pre-norm architecture.

**Core Methods:**

Table 21.4 lists the methods on this class.

Table 21.4: Core methods on the TransformerBlock class.

| Method                  | Signature                                                         | Description                                           |
|-------------------------|-------------------------------------------------------------------|-------------------------------------------------------|
| <code>forward</code>    | <code>forward(x: Tensor, mask: Tensor = None) -&gt; Tensor</code> | Process sequence through attention and MLP sub-layers |
| <code>parameters</code> | <code>parameters() -&gt; List[Tensor]</code>                      | Returns all parameters from attention, norms, and MLP |

### 21.4.5 GPT

```
GPT(vocab_size: int, embed_dim: int, num_layers: int, num_heads: int, max_seq_len:
    int = 1024) -> GPT
```

Complete GPT model for autoregressive language modeling with token embeddings, positional encoding, stacked transformer blocks, and generation capability. The architecture combines token and positional embeddings, processes through multiple transformer blocks with causal masking, applies final layer normalization, and projects to vocabulary logits.

#### Core Methods:

Table 21.5 lists the methods on this class.

Table 21.5: Core methods on the complete GPT model class.

| Method                           | Signature                                                                                                     | Description                                                                 |
|----------------------------------|---------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------|
| <code>forward</code>             | <code>forward(tokens: Tensor) -&gt; Tensor</code>                                                             | Compute vocabulary logits for each position with causal masking             |
| <code>generate</code>            | <code>generate(prompt_tokens: Tensor, max_new_tokens: int = 50, temperature: float = 1.0) -&gt; Tensor</code> | Autoregressively generate text using temperature-controlled sampling        |
| <code>parameters</code>          | <code>parameters() -&gt; List[Tensor]</code>                                                                  | Returns all model parameters from embeddings, blocks, and output head       |
| <code>_create_causal_mask</code> | <code>_create_causal_mask(seq_len: int) -&gt; Tensor</code>                                                   | Internal method creating upper triangular mask for autoregressive attention |

## 21.5 Core Concepts

This section explores the architectural innovations that make transformers the dominant deep learning architecture. Understanding these concepts deeply will prepare you for both implementing transformers and designing novel architectures.

### 21.5.1 Layer Normalization: The Stability Foundation

Without normalization, training a network with dozens of layers becomes nearly impossible: activation distributions drift between layers and gradients explode or vanish. Layer norm pins the distribution back to zero mean and unit variance at every step.

Unlike batch normalization, which mixes statistics across the batch dimension, layer norm normalizes each sample independently across its features. That independence is what makes it work for variable-length sequences: a batch containing a 10-token tweet and a 500-token paragraph would give batch norm meaningless mixed statistics, but layer norm treats each position on its own.

Here's the complete implementation, showing how a few targeted statistical operations keep the network's forward and backward passes perfectly scaled:

The code in `?@lst-13-layernorm` makes this concrete.

```

class LayerNorm:
    def __init__(self, normalized_shape, eps=1e-5):
        self.normalized_shape = normalized_shape
        self.eps = eps

        # Learnable parameters initialized to identity transform
        self.gamma = Tensor(np.ones(normalized_shape), requires_grad=True)
        self.beta = Tensor(np.zeros(normalized_shape), requires_grad=True)

    def forward(self, x):
        # Compute statistics across last dimension (features)
        mean = x.mean(axis=-1, keepdims=True)
        diff = x - mean
        variance = (diff * diff).mean(axis=-1, keepdims=True)

        # Normalize to zero mean, unit variance
        std = Tensor(np.sqrt(variance.data + self.eps))
        normalized = (x - mean) / std

        # Apply learnable transformation
        return normalized * self.gamma + self.beta

```

: Listing 13.1 — LayerNorm implementation. Computes per-sample mean and variance over the feature axis, then applies a learnable affine transform. {#lst-13-layernorm}

The formula is small:  $\text{output} = (x - \mu) / \sigma * \gamma + \beta$ . The effect is large. Forcing every activation back to a consistent scale keeps gradients in a sane range, which is the difference between a 12-layer model that trains and a 12-layer model that diverges in the first 100 steps. The learnable `gamma` and `beta` let the model recover any distribution it needs, so the normalization step costs you no expressiveness.

The `eps = 1e-5` guards against the rare case where every feature in a row is identical and variance hits zero. Without it, you divide by zero on the first batch.

### 21.5.2 Pre-Norm Architecture and Residual Connections

The original transformer (Vaswani et al., 2017) put layer norm *after* each sub-layer. Modern transformers put it *before*. The fix sounds trivial — swap the order — but it is what lets you train 24, 48, or 100 layers without warmup tricks. The pattern is: normalize, transform, add residual.

Residual connections are the gradient highways. When you write  $x + f(x)$ , backpropagation gets two paths home: through the transformation  $f$ , and straight through the  $+ x$  shortcut. Even if  $\partial f / \partial x$  is tiny, the shortcut contributes a clean 1 to the gradient, so signal still reaches the early layers in a 100-layer stack.

Here’s how the transformer block implements pre-norm with residuals:

The code in `?@lst-13-pre-norm-forward` makes this concrete.

```

def forward(self, x, mask=None):
    # First sub-layer: attention with pre-norm
    normed1 = self.ln1.forward(x)
    attention_out = self.attention.forward(normed1, mask)
    x = x + attention_out # Residual connection

    # Second sub-layer: MLP with pre-norm

```

```

normed2 = self.ln2.forward(x)
mlp_out = self.mlp.forward(normed2)
output = x + mlp_out  # Residual connection

return output

```

: **Listing 13.2 — Pre-norm transformer block forward pass.** LayerNorm runs before each sub-layer while residual additions write back into the unnormalized stream. {#lst-13-pre-norm-forward}

Notice the asymmetry: each sub-layer *reads* a normalized copy of the input but *writes* its contribution back into the unnormalized residual stream. The normalized path keeps the sub-layer well-behaved; the residual path keeps information flowing intact across the depth of the network.

### 21.5.3 The MLP: Computational Capacity Through Expansion

Attention handles relationships *between* tokens. The MLP handles transformation *within* each token, position by position, independently. Together they cover both axes of the sequence.

The standard recipe is wide-then-narrow: expand the embedding to  $4\times$  its width, apply GELU, project back. The expansion gives the model a high-dimensional scratch space to disentangle features; the projection forces it to compress what matters back into the residual stream. The  $4\times$  ratio is empirical, not principled — it is what worked in the original paper and has stuck ever since.

The code in ?@lst-13-mlp makes this concrete.

```

class MLP:
    def __init__(self, embed_dim, hidden_dim=None):
        if hidden_dim is None:
            hidden_dim = 4 * embed_dim  # Standard 4x expansion

        self.linear1 = Linear(embed_dim, hidden_dim)
        self.gelu = GELU()
        self.linear2 = Linear(hidden_dim, embed_dim)

    def forward(self, x):
        hidden = self.linear1.forward(x)
        hidden = self.gelu.forward(hidden)
        output = self.linear2.forward(hidden)
        return output

```

: **Listing 13.3 — Position-wise MLP.** Expands to  $4\times$  hidden width, applies GELU, and projects back. Holds most of the block’s parameters. {#lst-13-mlp}

GELU replaced ReLU in transformer models because it gates smoothly instead of with a hard cutoff at zero, which gives cleaner gradients for language modeling. The choice matters less than the width: most modern variants (GELU, SwiGLU, GeGLU) train to similar loss curves at this scale.

The MLP dominates the parameter count. For  $\text{embed\_dim} = 512$ , the first projection has  $512 \times 2048 + 2048 = 1,050,624$  (~1.05M) parameters, the second has  $2048 \times 512 + 512 = 1,049,088$  (~1.05M), for ~2.1M per block. In a 12-layer model that is ~25.2M parameters from MLPs alone — more than attention and embeddings combined.

### 21.5.4 Causal Masking for Autoregressive Generation

GPT is an autoregressive model: it predicts each token based only on previous tokens. During training, the model sees the entire sequence, but causal masking ensures position  $i$  cannot attend to positions  $j > i$ . This prevents information leakage from the future.

The causal mask is an upper triangular matrix filled with negative infinity:

```
def create_causal_mask(seq_len: int) -> Tensor:
    # Lower triangle = 1 (can attend), upper triangle = 0 (cannot attend)
    mask = np.tril(np.ones((seq_len, seq_len), dtype=np.float32))
    return Tensor(mask[np.newaxis, :, :])
```

For a 4-token sequence, this creates:

```
[[1, 0, 0, 0], # Position 0 only sees itself
 [1, 1, 0, 0], # Position 1 sees 0, 1
 [1, 1, 1, 0], # Position 2 sees 0, 1, 2
 [1, 1, 1, 1]] # Position 3 sees everything
```

Inside attention, these zeros become  $-\text{inf}$  in the logits before softmax. After softmax,  $-\text{inf}$  collapses to exactly 0 probability, so future positions contribute nothing to the weighted sum. The mask is what lets you train on a 2048-token sequence in a single parallel pass while still computing every prediction as if you only knew the past.

### 21.5.5 Complete Transformer Block Architecture

The transformer block is where all components unite into a coherent processing unit. Each block transforms the input sequence through two sub-layers: multi-head self-attention and MLP, each wrapped with layer normalization and residual connections.

The code in `lst-13-transformer-block` makes this concrete.

```
class TransformerBlock:
    def __init__(self, embed_dim, num_heads, mlp_ratio=4):
        self.attention = MultiHeadAttention(embed_dim, num_heads)
        self.ln1 = LayerNorm(embed_dim) # Before attention
        self.ln2 = LayerNorm(embed_dim) # Before MLP
        hidden_dim = int(embed_dim * mlp_ratio)
        self.mlp = MLP(embed_dim, hidden_dim)

    def forward(self, x, mask=None):
        # First sub-layer: attention with residual
        normed1 = self.ln1.forward(x)
        attention_out = self.attention.forward(normed1, mask)
        x = x + attention_out # Residual connection

        # Second sub-layer: MLP with residual
        normed2 = self.ln2.forward(x)
        mlp_out = self.mlp.forward(normed2)
        output = x + mlp_out # Residual connection

    return output
```



: **Listing 13.4 — Complete TransformerBlock class.** Wires MultiHeadAttention, two LayerNorms, and an MLP into a single pre-norm block with residual connections. {#lst-13-transformer-block}

Think of the data flow as a *residual stream*: the input embeddings enter, and every sub-layer adds its contribution on top without overwriting what came before. By the final block, the stream is the original embeddings plus contributions from every attention and MLP sub-layer in the stack — like a stack of transparencies, each adding detail to the same underlying image.

This is why transformers scale to hundreds of layers while plain MLPs choke at ten. Each layer’s job is to *adjust* the stream, not replace it. Backprop pushes gradients through these small additive corrections, and the residual shortcuts keep them from decaying as they travel.

### 21.5.6 Parameter Scaling and Memory Requirements

Two scaling laws govern transformer cost. Parameters scale roughly with  $\text{embed\_dim}^2$  (because attention and MLP weights are square-ish matrices in that dimension). Attention activations scale with  $\text{seq\_len}^2$  (every token attends to every other). Either one can dominate your hardware budget; both will, eventually.

For a single transformer block with  $\text{embed\_dim} = 512$  and  $\text{num\_heads} = 8$ :

Table 21.6 breaks down the parameter count for a single transformer block.

Table 21.6: Parameter count breakdown for a single 512-dim, 8-head transformer block.

| Component              | Parameters   | Calculation                                            |
|------------------------|--------------|--------------------------------------------------------|
| Multi-Head Attention   | ~1.05M       | $4 \times (512 \times 512)$ for Q, K, V, O projections |
| Layer Norm 1           | 1K           | $2 \times 512$ for gamma, beta                         |
| MLP                    | ~2.1M        | $(512 \times 2048 + 2048) + (2048 \times 512 + 512)$   |
| Layer Norm 2           | 1K           | $2 \times 512$ for gamma, beta                         |
| <b>Total per block</b> | <b>~3.2M</b> | Dominated by MLP, then attention                       |

For a complete GPT model, add embeddings and output projection:

- **Token embeddings:**  $\text{vocab\_size} \times \text{embed\_dim} = 50000 \times 512 = 25.6\text{M}$
- **Position embeddings:**  $\text{max\_seq\_len} \times \text{embed\_dim} = 2048 \times 512 = 1.0\text{M}$
- **Transformer blocks:**  $12 \times 3.2\text{M} = 37.8\text{M}$
- **Output projection:**  $\text{embed\_dim} \times \text{vocab\_size}$  (typically tied to token embeddings, so 0 extra)

Grand total: **~64M parameters** for this configuration. GPT-2 small was ~117M; GPT-3 is 175B. The arithmetic does not change — you just multiply the same blocks by larger dimensions and more layers.

Memory at training time has three components:

1. **Parameter memory** — linear with model size, stored once
2. **Activation memory** — needed for backprop, grows with batch size and sequence length
3. **Attention memory** — quadratic in sequence length, the bottleneck that bites first

The attention bottleneck is why long context is expensive. For a batch of 4 sequences with 8 heads in float32:

Table 21.7 shows how attention memory grows with sequence length.

Table 21.7: Attention memory growth with sequence length for a fixed-size transformer.

| Sequence Length | Attention Matrix Size                | Memory (MB) |
|-----------------|--------------------------------------|-------------|
| 512             | $4 \times 8 \times 512 \times 512$   | 32.0        |
| 1024            | $4 \times 8 \times 1024 \times 1024$ | 128.0       |
| 2048            | $4 \times 8 \times 2048 \times 2048$ | 512.0       |
| 4096            | $4 \times 8 \times 4096 \times 4096$ | 2048.0      |

Double the sequence length, quadruple the memory — *per layer*. A 12-layer model at 4K context burns 24 GB just on attention matrices. This is the wall that drove FlashAttention, sparse attention, and linear attention into existence.

Per transformer block, the cost profile decomposes cleanly into two regimes:

- **Attention:** compute is  $O(N^2 \cdot d)$ , memory is  $O(N^2)$  — quadratic in sequence length  $N$ .
- **MLP:** compute is  $O(N \cdot d^2)$ , memory is  $O(N \cdot d)$  — linear in  $N$  but quadratic in the hidden width  $d$ .

Attention overtakes the MLP the moment  $N > d$ , which for every production LLM (where  $N$  is thousands of tokens and  $d$  is a few hundred to a few thousand) is always. That crossover is why the attention matrix — not the weight matrices — is the first thing that blows up your HBM budget.

### **i** Systems Implication: KV-Cache Memory Growth During Autoregressive Decoding

The  $O(N^2)$  scaling story above is about *training*. Autoregressive *inference* has a different — and equally brutal — memory problem: the **KV cache**. When generating tokens one at a time, naive attention would recompute the keys and values for every previous token at every step (the  $66\times$  overhead shown in the generation efficiency question below). Production inference instead caches  $k$  and  $v$  for every token ever seen, trading compute for memory.

The cache size per request is  $2 \times \text{num\_layers} \times \text{num\_heads} \times \text{head\_dim} \times \text{seq\_len} \times \text{bytes}$  — two copies (K and V), one per layer, per head. For a 70 B Llama-style model at 32K context in FP16, that is roughly **20 GB per request**, often larger than the model's own weight footprint. At scale this is the dominant cost: inference servers are memory-bound on the KV cache, not FLOP-bound on the matmuls. This is why recent architectures use **multi-query attention** (MQA) and **grouped-query attention** (GQA) — they share keys and values across heads to shrink the cache by  $4\text{--}8\times$ , trading a small amount of model quality for a huge cut in serving memory. Every optimization in modern LLM inference stacks — paged attention, prefix caching, speculative decoding — is ultimately a strategy for managing the KV cache.

## 21.6 Production Context

### 21.6.1 Your Implementation vs. PyTorch

Your transformer implementation and PyTorch's production transformers share the same architectural principles. The differences lie in optimization: PyTorch uses fused CUDA kernels, memory-efficient attention, and various tricks for speed and scale.

Table 21.8 places your implementation side by side with the production reference for direct comparison.

Table 21.8: Feature comparison between TinyTorch transformers and production PyTorch.

| Feature             | Your Implementation         | PyTorch                                 |
|---------------------|-----------------------------|-----------------------------------------|
| <b>Architecture</b> | Pre-norm transformer blocks | Pre-norm (modern) or post-norm (legacy) |
| <b>Attention</b>    | Standard scaled dot-product | FlashAttention, sparse attention        |
| <b>Memory</b>       | Full attention matrices     | KV caching, memory-efficient attention  |
| <b>Precision</b>    | Float32                     | Mixed precision (FP16/BF16)             |
| <b>Parallelism</b>  | Single device               | Model parallel, pipeline parallel       |
| <b>Efficiency</b>   | Educational clarity         | Production optimization                 |

### 21.6.2 Code Comparison

The following comparison shows equivalent transformer usage in TinyTorch and PyTorch. The API patterns are nearly identical because your implementation follows production design principles.

## 21.7 Your TinyTorch

```
from tinytorch.core.transformers import TransformerBlock, GPT

# Create transformer block
block = TransformerBlock(embed_dim=512, num_heads=8)
output = block.forward(x)

# Create complete GPT model
model = GPT(vocab_size=50000, embed_dim=768, num_layers=12, num_heads=12)
logits = model.forward(tokens)
generated = model.generate(prompt, max_new_tokens=50, temperature=0.8)
```

## 21.8 PyTorch

```
import torch.nn as nn

# PyTorch transformer block (using nn.TransformerEncoderLayer)
block = nn.TransformerEncoderLayer(d_model=512, nhead=8, dim_feedforward=2048)
output = block(x)

# Complete model (using HuggingFace transformers)
from transformers import GPT2LMHeadModel, GPT2Tokenizer
```

```
model = GPT2LMHeadModel.from_pretrained("gpt2")
outputs = model.generate(input_ids, max_new_tokens=50, temperature=0.8)
```

Let's walk through the key similarities and differences:

- **Line 1-2 (Block creation):** Both create transformer blocks with identical parameters. PyTorch uses `TransformerEncoderLayer` while you built `TransformerBlock` from scratch.
- **Line 3 (Forward pass):** Both process sequences with identical semantics. Your implementation explicitly shows attention and MLP; PyTorch's is identical internally.
- **Line 5-6 (Model creation):** Both create complete language models. PyTorch typically uses pre-trained models via HuggingFace; you build from scratch.
- **Line 7 (Generation):** Both support autoregressive generation with temperature control. PyTorch adds beam search, top-k/top-p sampling, and other advanced techniques.

#### 💡 What's Identical

The core architecture, pre-norm pattern, residual connections, and causal masking are identical. When you debug transformer models in PyTorch, you'll understand exactly what's happening because you built it yourself.

### 21.8.1 Why Transformers Matter at Scale

The architecture you just built is the same one that runs at the frontier. The difference is the dimensions.

- **GPT-3 (175B parameters)** — 350 GB just to store weights in float16, 700 GB for mixed-precision training state.
- **Training cost** — roughly \$4.6M in compute, ~10,000 GPUs for weeks.
- **Inference latency** — 100–200 ms to process 2048 tokens on optimized hardware.
- **Context scaling** — going from 2K to 32K context costs  $256\times$  more attention memory *per layer*.

Those numbers exist because the architecture you built has no free lunch baked in: every doubling of context squares the attention cost, every doubling of width squares the parameter count. The next chapter is where you start measuring those costs in your own model — and the chapters after that are where you start beating them down.

#### 💡 Check Your Understanding — Transformers

Before moving on, verify you can articulate each of the following:

- Why the residual connection around attention+FFN ( $x + f(x)$ ) is critical for gradient flow through deep stacks — the  $+1$  in  $\partial y / \partial x = 1 + \partial f / \partial x$  is a shortcut that lets signal reach early layers even when the sub-layer Jacobian is tiny.
- How **pre-norm** (LayerNorm *before* each sub-layer) differs from the original *post-norm* transformer, and why it's what lets modern models train 24, 48, or 100+ layers without learning-rate warmup tricks.
- Why the attention block's  $O(N^2)$  memory cost dominates over the MLP's  $O(N \cdot d^2)$  the moment  $N > d$  — and what that means for context-window choices in every production LLM.
- How the **KV cache** turns autoregressive decoding from  $O(N^2)$  recomputation into  $O(N)$  per step, why the cache itself can exceed the model's weight footprint, and why MQA/GQA/paged attention exist to shrink it.

If any of these feels fuzzy, revisit the *Core Concepts* section (especially *Pre-Norm Architecture and Residual Connections* and *Parameter Scaling and Memory Requirements*) before moving on.

## 21.9 Self-Check Questions

Test your understanding of transformer architecture and scaling with these systems thinking questions.

### Q1: Attention Memory Calculation

A transformer with `batch_size=8`, `num_heads=16`, `seq_len=2048` computes attention matrices at each layer. How much memory does one layer's attention matrices consume? How does this scale if you double the sequence length to 4096?

#### 💡 Answer

**Attention matrix size:**  $\text{batch\_size} \times \text{num\_heads} \times \text{seq\_len} \times \text{seq\_len} = 8 \times 16 \times 2048 \times 2048 = 536,870,912$  elements

**Memory:**  $536,870,912 \times 4 \text{ bytes (float32)} = 2,147,483,648 \text{ bytes} = \mathbf{2.0 \text{ GB}}$

**Doubling sequence length to 4096:**  $= 8 \times 16 \times 4096 \times 4096 = 2,147,483,648$  elements = **8.0 GB**

**Scaling:** doubling sequence length quadruples memory (4× increase). This quadratic scaling is why long context is expensive — and why sparse attention, FlashAttention, and linear-attention variants exist.

### Q2: Parameter Distribution Analysis

For a GPT model with `vocab_size=50000`, `embed_dim=768`, `num_layers=12`, `num_heads=12`, calculate approximate total parameters. Which component dominates the parameter count: embeddings or transformer blocks?

#### 💡 Answer

**Token Embeddings:**  $50000 \times 768 = 38.4\text{M}$

**Position Embeddings:**  $2048 \times 768 = 1.6\text{M}$  (assuming `max_seq_len=2048`)

**Transformer Blocks:** each block has ~7.1M parameters with `embed_dim=768`

- Attention:  $4 \times (768 \times 768) = \sim 2.4\text{M}$
- MLP:  $(768 \times 3072 + 3072) + (3072 \times 768 + 768) = \sim 4.7\text{M}$
- Layer norms: negligible
- Per block: ~7.1M
- Total blocks:  $12 \times 7.1\text{M} = 85\text{M}$

**Output Projection:** usually tied to embeddings (0 additional)

**Total:**  $38.4\text{M} + 1.6\text{M} + 85\text{M} \approx \mathbf{125\text{M parameters}}$  — close to GPT-2 small.

**Dominant component:** transformer blocks (85M) outweigh embeddings (40M). As models scale, blocks pull further ahead because they grow with  $\text{embed\_dim}^2$  while embeddings grow only linearly with vocab size.

### Q3: Residual Connection Benefits

Why do transformers use residual connections ( $x + f(x)$ ) rather than just  $f(x)$ ? How do residual connections affect gradient flow during backpropagation in a 24-layer transformer?

#### 💡 Answer

**Without residual connections** ( $y = f(x)$ ):

- Gradients must flow through every transformation layer.
- Each layer's Jacobian can shrink gradients (vanishing) or amplify them (exploding).
- Across 24 layers, gradients can collapse to  $\sim 0$  or blow up to  $\sim \infty$ .

**With residual connections** ( $y = x + f(x)$ ):

- Backprop gives  $\partial y / \partial x = 1 + \partial f / \partial x$ .
- The  $+1$  is a direct gradient path that does not depend on the sub-layer's weights.
- Even if  $\partial f / \partial x$  is small, the  $+1$  keeps the signal alive.
- The result is a stack of “gradient highways” running the depth of the network.

**24-layer impact:** without residuals, a 0.9 per-layer attenuation compounds to  $0.9^{24} \approx 0.08$ . With residuals, the  $+1$  shortcut delivers gradients to early layers at full strength. This is why transformers scale to 100+ layers while plain feed-forward nets struggle past 10.

#### Q4: Autoregressive Generation Efficiency

Your `generate()` method processes the entire sequence for each new token. For generating 100 tokens with prompt length 50, how many total forward passes occur? Why is this inefficient?

#### 💡 Answer

**Current implementation:** for each of 100 new tokens, reprocess the entire sequence.

- Token 1: process 50 tokens (prompt)
- Token 2: process 51 tokens (prompt + 1)
- Token 3: process 52 tokens
- ...
- Token 100: process 149 tokens

**Total forward passes:**  $50 + 51 + 52 + \dots + 149 = 9,950$  token processings

**Why inefficient:** attention recomputes key/value projections for every previous token at every step, even though those projections do not change. The key/value for position 50 is recomputed 100 times.

**KV Caching optimization:** store computed key/value projections for previous tokens.

- Each new token only computes its own key/value.
- Attention uses cached keys/values from previous tokens.
- Total computation:  $50$  (prompt) +  $100$  (new tokens) = **150 token processings**.

**Speedup:**  $9950 / 150 \approx 66\times$  faster for this example. The ratio grows with generation length, which is why KV caching is non-negotiable in production inference stacks.

#### Q5: Layer Normalization vs Batch Normalization

Why do transformers use layer normalization instead of batch normalization? Consider a batch with sequences of varying lengths: [10 tokens, 50 tokens, 100 tokens].

#### 💡 Answer

**Batch Normalization** normalizes across the batch dimension:

- For position 5, statistics would mix all three sequences.
- But sequence 1 has no position 50, sequence 2 has no position 100.
- With padding, statistics get contaminated by pad tokens.
- Depends on batch composition: different batches give different statistics.

**Layer Normalization** normalizes across features for each sample:

- Each position is normalized independently:  $(x - \text{mean}(x)) / \text{std}(x)$ .
- Position 5 of sequence 1 has no influence on position 50 of sequence 2.
- No dependency on batch composition.
- Works naturally with variable-length sequences.

**Example:** for a tensor of shape (batch=3, seq=10, features=768):

- Batch norm computes  $10 \times 768$  statistics across the batch dimension (problematic).
- Layer norm computes  $3 \times 10$  statistics across the feature dimension (independent).

**Why it matters:** transformers process variable-length sequences. Layer norm treats each position on its own, which makes it robust to length variation and batch composition. Batch norm depends on its batch; layer norm does not.

## 21.10 Key Takeaways

- **A transformer block is a recipe, not a mystery:** pre-norm  $\rightarrow$  attention  $\rightarrow$  residual, pre-norm  $\rightarrow$  MLP  $\rightarrow$  residual. Stack twelve of those between embeddings and a language head and you have GPT.
- **Residuals are the gradient highway:**  $x + f(x)$  gives backprop a +1 shortcut that keeps signal alive through arbitrarily deep stacks — without them, a 24-layer network's gradients decay to zero.
- **Attention's  $O(N^2)$  memory is the first wall you hit:** doubling sequence length quadruples attention memory per layer; long context is expensive because of geometry, not implementation.
- **The KV cache dominates inference memory:** at serving time,  $2 \times \text{layers} \times \text{heads} \times \text{head\_dim} \times \text{seq\_len} \times \text{bytes}$  often eclipses the model weights themselves, which is why MQA, GQA, and paged attention now ship in every production LLM stack.

**Coming next:** Module 14 opens the Optimization Tier with measurement — profiling your transformer's forward pass to see exactly where time and memory are spent before you try to cut either.

## 21.11 Further Reading

The transition from a theoretical transformer diagram to a planet-scale language model is governed by ruthless hardware constraints and brilliant systems engineering. For students seeking to understand the architectural turning points and the hardware-software co-design that made massive models possible, study these foundational texts.

### 21.11.1 Seminal Papers

- **Attention Is All You Need** - Vaswani et al. (2017). The paper that introduced the transformer architecture, revolutionizing sequence modeling. Describes multi-head attention, positional encoding, and the encoder-decoder structure. **Systems Implication:** Eradicated the sequential compute bottleneck inherent to RNNs. By processing entire sequences simultaneously, the transformer perfectly aligns with the SIMD architecture of modern GPUs, enabling massive parallelization and achieving near-peak hardware utilization. [arXiv:1706.03762](#)
- **Language Models are Few-Shot Learners (GPT-3)** - Brown et al. (2020). Demonstrates scaling laws and emergent capabilities of large language models. Shows how transformer performance improves predictably with immense scale. **Systems Implication:** At 175 billion parameters, the model shattered the memory capacity of a single GPU. This forced the systems engineering community to rapidly mature and deploy complex 3D parallelism strategies (data, tensor, and pipeline parallelism) across clusters of thousands of distinct compute nodes. [arXiv:2005.14165](#)
- **FlashAttention: Fast and Memory-Efficient Exact Attention** - Dao et al. (2022). Overcomes the  $O(n^2)$  memory scaling limit of attention using IO-aware algorithms, enabling practical long-context processing. **Systems Implication:** Leveraged precise, hardware-aware SRAM tiling to completely bypass the materialization of the gigantic quadratic attention matrix in HBM. This critical optimization shifted attention from being a catastrophic memory-bound operation back to a highly efficient compute-bound workload. [arXiv:2205.14135](#)



- **On Layer Normalization in the Transformer Architecture** - Xiong et al. (2020). Mathematically analyzes pre-norm versus post-norm architectures, revealing why pre-norm enables the training of much deeper networks. **Systems Implication:** Pre-norm architectures prevent catastrophic gradient vanishing or exploding during backpropagation across hundreds of layers. By guaranteeing gradient stability, these architectures ensure that massive, distributed training runs remain convergent, saving millions of dollars in idle or wasted GPU compute. [arXiv:2002.04745](#)

### 21.11.2 Additional Resources

- **Blog post:** [“The Illustrated Transformer”](#) by Jay Alammar - Visual walkthrough of transformer architecture with clear diagrams
- **Paper:** [“Scaling Laws for Neural Language Models”](#) - Kaplan et al. (2020) - Mathematical analysis of how performance scales with parameters, data, and compute
- **Implementation:** [HuggingFace Transformers library](#) - Production transformer implementations to compare with your code

## 21.12 What’s Next

You finished the Architecture Tier. You have a transformer that trains, generates text, and matches the structural blueprint of GPT. Everything from here is about making it *fast, small, and deployable*.

Before that, the next two chapters take the Architecture Tier on a historical test drive. The **Architecture Milestones** — a LeNet-style CNN on TinyDigits, with optional CIFAR-10 scale-up, and the 2017 Transformer attention test — run your `Conv2d`, `MaxPool2d`, multi-head attention, and `TransformerBlock` on problems those architectures were built to solve. Convolutional networks exploit spatial structure; attention clears sequence-reversal tasks RNNs can’t. Both are proof that the layers you just wrote behave the way the original papers claimed.

### **i** Coming Up: Architecture Milestones, then Module 14 — Profiling

First: two Architecture Milestones exercise your Conv/Pool layers (TinyDigits, with optional CIFAR-10 scale-up) and your attention stack (sequence tasks) on their landmark problems. Then the Optimization Tier opens with the only honest place to start: measurement. Before you optimize anything, you need to know *where the time goes*. In Module 14 you instrument your transformer’s forward pass and answer concrete questions — how much of a step is spent in attention versus the MLP, how memory grows with sequence length on your machine, and which layer is the actual bottleneck. Every optimization in the chapters that follow (quantization, kernel fusion, KV caching) targets a number you measured in Module 14.

### Where your transformer goes from here:

Table 21.9 traces how this module is reused by later parts of the curriculum.

Table 21.9: **How transformers feed into profiling, quantization, and capstone modules.**

| Module                  | What It Does                      | Your Transformer In Action                                                       |
|-------------------------|-----------------------------------|----------------------------------------------------------------------------------|
| <b>14: Profiling</b>    | Measure performance bottlenecks   | <code>profiler.analyze(model.forward(x))</code> reveals where time and memory go |
| <b>15: Quantization</b> | Reduce precision to 8-bit / 4-bit | Shrink the 64M model to a fraction of its size with minimal accuracy loss        |



| Module       | What It Does          | Your Transformer In Action                                   |
|--------------|-----------------------|--------------------------------------------------------------|
| 20: Capstone | Production deployment | Serve your transformer end-to-end with the systems you built |

## 21.13 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



---

PART IV

# Architecture Milestones

---



## 🔥 Chapter 22

# Milestone 04: The CNN Revolution (1998)

## i Milestone Info

**Architecture Milestone** | Difficulty: ●●●○ | Time: 1–2 hours (incl. training) | Prerequisites: Modules 01–09

## 💡 What You'll Learn

- Why spatial structure matters:  $100\times$  fewer parameters, 50% better accuracy
- How weight sharing enables translation invariance
- The hierarchical feature learning that powers all computer vision

## 22.1 Overview

This is the first **Architecture Milestone**. The Foundation Milestones proved your training loop learns; the next two prove your *architectures* do the work they were invented for. Here you pick up the `Conv2d` and `MaxPool2d` layers you just built in Module 09 and first validate them on `TinyDigits`, then optionally scale to natural images.

1998. Yann LeCun deploys LeNet-5: a convolutional neural network that reads handwritten zip codes for the US Postal Service and dollar amounts on bank checks for NCR. It is not a research demo. It is production software, sorting mail and clearing checks at industrial scale — the first commercial success of deep learning.

The breakthrough is structural. Images are not bags of pixels; nearby pixels matter more than distant ones, and the same edge detector works whether you put it in the corner or the center. Exploit those two facts — local connectivity and weight sharing — and a convolutional network matches a dense one's accuracy with two orders of magnitude fewer parameters.

You are about to reproduce those same principles using your own `Conv2d` and `MaxPool2d` from Module 09. The default milestone runs offline on `TinyDigits`; the `CIFAR-10` script is the optional scale-up path when you want the natural-image benchmark.

## 22.2 What You'll Build

CNNs that exploit image structure:

1. **TinyDigits** — prove convolution beats MLPs on  $8\times 8$  images
2. **CIFAR-10** — scale to natural color images ( $32\times 32$  RGB)

Images --> Conv --> ReLU --> Pool --> Conv --> ReLU --> Pool --> Flatten --> Linear -->  
Classes

## 22.3 Prerequisites

Table 22.1 lists the modules you need to have completed before starting.

Table 22.1: Prerequisite modules for the CNN milestone.

| Module | Component             | What It Provides               |
|--------|-----------------------|--------------------------------|
| 01–08  | Foundation + Training | Complete training pipeline     |
| 09     | <b>Convolutions</b>   | <b>Your Conv2d + MaxPool2d</b> |

## 22.4 Running the Milestone

Before running, ensure you have completed Modules 01–09. You can check your progress:

```
tito module status
```

```
cd milestones/04_1998_cnn
```

```
# Default: TinyDigits (works offline)
```

```
tito milestone run 04
```

```
# Expected: ~90% accuracy (vs ~80% MLP)
```

```
# Or run the script directly
```

```
cd milestones/04_1998_cnn
```

```
python 01_lecun_tinydigits.py
```

```
# Optional scale-up: CIFAR-10 (requires download)
```

```
python 02_lecun_cifar10.py
```

```
# Expected: 65–75% accuracy
```

## 22.5 Expected Results

Table 22.2 records the accuracy and runtime you should expect to see.

Table 22.2: Expected accuracy for the CNN milestone on TinyDigits and CIFAR-10.

| Script          | Dataset              | Architecture | Accuracy | vs MLP             |
|-----------------|----------------------|--------------|----------|--------------------|
| 01 (TinyDigits) | 1K train, 8×8        | Simple CNN   | ~90%     | +10% improvement   |
| 02 (CIFAR-10)   | 50K train, 32×32 RGB | Deeper CNN   | 65–75%   | MLPs struggle here |

## 22.6 The Aha Moment: Structure Matches Reality

An MLP sees an image as 3,072 unrelated numbers. It does not know that pixel (0,0) is next to pixel (0,1). It learns brittle correlations like “if pixel 1,234 is bright and pixel 2,891 is dark...” — patterns tied to absolute positions, which fall apart the moment the cat shifts a few pixels to the left.

A CNN bakes spatial structure into the architecture itself:

1. **Local connectivity** — each neuron only looks at a small neighborhood ( $3 \times 3$  or  $5 \times 5$ ). Edges, corners, and textures are local patterns; the network does not need a global view to detect them.
2. **Weight sharing** — one filter scans the entire image. “Cat in the top-left” and “cat in the bottom-right” trigger the same feature detector, so the network learns the concept once instead of 1,024 times.
3. **Translation invariance** — pooling makes the output insensitive to small shifts. The network learns *that* a cat is present, not *where* the pixels happened to land.

The optional scale-up numbers:

- MLP on CIFAR-10: ~100M parameters, ~50% accuracy (memorizes pixel positions, fails to generalize)
- Your CNN: ~1M parameters, 75%+ accuracy (learns reusable features that compose)

**100× fewer parameters. 25 percentage points more accuracy.** That is what happens when the architecture matches the data.

The default milestone validates your implementations on TinyDigits. The optional Part 2 scales them: 50,000 natural color images,  $32 \times 32 \times 3 = 3,072$  dimensions per image, 10 categories (airplanes, cars, birds, cats, ships...). This is the hard problem.

Your `DataLoader` streams batches from disk. Your `Conv2d` layers extract features hierarchically — first layer finds edges, second finds textures, third finds object parts. Your `MaxPool2d` shrinks the spatial map while preserving what matters.

When the optional CIFAR-10 run prints a strong test accuracy, sit with it for a second. You did not download a pretrained model. Every tensor op, every gradient, every parameter update is code you wrote — running on your laptop, training a model that would have made headlines twenty years ago. That is systems engineering.

## 22.7 Your Code Powers This

Table 22.3 names the TinyTorch components that power this milestone.

Table 22.3: TinyTorch components that power the CNN milestone.

| Component                     | Your Module | What It Does                   |
|-------------------------------|-------------|--------------------------------|
| <code>Tensor</code>           | Module 01   | Stores images and feature maps |
| <code>Conv2d</code>           | Module 09   | Your convolutional layers      |
| <code>MaxPool2d</code>        | Module 09   | Your pooling layers            |
| <code>ReLU</code>             | Module 02   | Your activation functions      |
| <code>Linear</code>           | Module 03   | Your classifier head           |
| <code>CrossEntropyLoss</code> | Module 04   | Your loss computation          |
| <code>DataLoader</code>       | Module 05   | Your batching pipeline         |
| <code>backward()</code>       | Module 06   | Your autograd engine           |

## 22.8 Historical Context

LeNet-5 was deployed for zip code recognition at the US Postal Service and check-amount reading at NCR — the first neural networks to ship in production at meaningful scale.

CIFAR-10 (2009) became the standard pre-ImageNet benchmark. Reaching 70%+ on it was the signal the field was ready for the next jump in scale.

The 2012 “ImageNet moment” — AlexNet — applied the same CNN principles to 1.2 million images on GPUs. The blueprint was already in LeCun’s 1998 paper. The hardware just had to catch up.

## 22.9 Systems Insights

- **Memory:** ~1M parameters (weight sharing dramatically reduces vs dense)
- **Compute:** Convolution is compute-intensive but highly parallelizable
- **Architecture:** Hierarchical feature learning (edges → textures → objects)

## 22.10 What’s Next

CNNs are the right inductive bias for grid-structured data, but most of the world’s interesting signals are sequential — text, audio, time series. **Milestone 05** introduces **Transformers**, the architecture that ate sequence modeling first and, eventually, vision itself.

## 22.11 Further Reading

- **LeNet Paper:** LeCun et al. (1998). “Gradient-based learning applied to document recognition”
- **CIFAR-10:** Krizhevsky (2009). “Learning Multiple Layers of Features from Tiny Images”
- **AlexNet:** Krizhevsky, A., Sutskever, I., & Hinton, G. (2012). “ImageNet Classification with Deep Convolutional Neural Networks”



## 🔥 Chapter 23

# Milestone 05: The Transformer Era (2017)

## i Milestone Info

**Architecture Milestone** | Difficulty: ●●●○ | Time: 30–60 min | Prerequisites: Modules 01–08, 11–13

## 💡 What You'll Learn

- Why attention beats sequential processing (direct access, no bottleneck)
- How transformers handle long-range dependencies that defeat RNNs
- The architecture that powers GPT, BERT, and every modern LLM

## 23.1 Overview

**Imagine compressing an entire book into a single number.**

That's what an RNN does. The input — a sentence, a paragraph, an entire novel — squeezes through one fixed-size hidden state. Information from the beginning fades as new tokens arrive. By token 1000, token 1 is nearly forgotten.

In 2017, Vaswani et al. asked: *what if we just... didn't?* Their paper “Attention Is All You Need” proved that attention alone — no recurrence, no convolution — could match the state of the art. No sequential bottleneck. No information loss. Any token can directly attend to any other token.

Nearly a decade later, every frontier LLM is a transformer: GPT-4, Claude, Gemini, Llama. The architecture you finished in Modules 12–13 is the architecture behind the AI boom. This milestone proves your version works on tasks RNNs cannot solve.

## 23.2 What You'll Build

1. **Attention proof** (`01_vaswani_attention.py`) — three synthetic tasks (sequence reversal, copying, mixed prefixes) that working self-attention can solve and broken attention cannot.
2. **Optional corpus** — TinyTalks in `datasets/tinytalks/` for Q&A-style transformer experiments, independent of `tito` milestone run 05.

Tokens --> Embeddings --> [Attention --> FFN] x N --> Output

## 23.3 Prerequisites

Table 23.1 lists the modules you need to have completed before starting.

Table 23.1: Prerequisite modules for the Transformer milestone.

| Module | Component             | What It Provides                     |
|--------|-----------------------|--------------------------------------|
| 01–08  | Foundation + Training | Tensor, Layers, DataLoader, Training |
| 11     | Embeddings            | YOUR token + positional embeddings   |
| 12     | Attention             | YOUR multi-head self-attention       |
| 13     | Transformers          | YOUR LayerNorm + TransformerBlock    |

## 23.4 Running the Milestone

Before running, ensure you have completed Modules 01–13. You can check your progress:

```
tito module status
```

```
tito milestone run 05
```

Or:

```
cd milestones/05_2017_transformer
python 01_vaswani_attention.py
```

### i Synthetic vs TinyTalks

The milestone entrypoint uses **synthetic** sequences so attention can be verified quickly with no extra files. **TinyTalks** remains available under `datasets/tinytalks/` for teaching and follow-on work.

## 23.5 Expected Results

Table 23.2 records the accuracy and runtime you should expect to see.

Table 23.2: Expected success criteria and runtime for the Transformer milestone.

| Script                  | Task                       | Success Criteria                           | Time    |
|-------------------------|----------------------------|--------------------------------------------|---------|
| 01_vaswani_attention.py | Reversal / copy<br>/ mixed | See script thresholds (~95% / ~95% / ~90%) | Minutes |

## 23.6 The Aha Moment: Direct Access Everywhere

**Sequence reversal** demands every output position read every input position — a stress test for cross-position attention. **Copying** forces an identity-like alignment pattern. **Mixed prefixes** mirror the context-conditioned behavior of decoder-style language models.

RNNs choke on these tasks; the dependencies are too long for a single hidden state. Attention makes them learnable:

```
RNN:      h[t]    = f(h[t-1], x[t])          # Sequential, lossy
Attention: out[i] = sum(attn[i, j] * v[j])    # Parallel, direct
```

## 23.7 Your Code Powers This

Table 23.3 names the TinyTorch components that power this milestone.

Table 23.3: TinyTorch components that power the Transformer milestone.

| Component           | Your Module | What It Does                        |
|---------------------|-------------|-------------------------------------|
| TokenEmbedding      | Module 11   | YOUR learned token representations  |
| PositionalEmbedding | Module 11   | YOUR position encodings             |
| MultiHeadAttention  | Module 12   | YOUR self-attention mechanism       |
| TransformerBlock    | Module 13   | YOUR attention + feedforward blocks |
| LayerNorm           | Module 13   | YOUR normalization layers           |

**No PyTorch. No HuggingFace. Just YOUR code.**

## 23.8 Historical Context

The 2017 paper trained a 65M-parameter encoder-decoder on English↔German translation. The same blocks — same attention, same residual stream, same LayerNorm — scaled to GPT-3’s 175B parameters in 2020 with no architectural change. Bigger model, more data, longer training. That scaling story is what made transformers a paradigm rather than a paper. This milestone uses small synthetic tasks for fast feedback; **TinyTalks** in `datasets/tinytalks/` provides Q&A-style character-level text when you want a real corpus.

## 23.9 Systems Insights

- **Memory:** The attention matrix is  $O(n^2)$ . At 8K tokens, that’s 64M scores per head per layer — the cost that drives every long-context optimization (FlashAttention, sliding windows, KV cache eviction).
- **Compute:** Embarrassingly parallel across positions. RNNs serialize on the hidden state; transformers fill a GPU.
- **Order:** Attention is permutation-invariant. Position embeddings are how the model learns “first” from “last” — shuffle them and the model loses all sense of sequence.

## 23.10 What’s Next

You can *build* a transformer. The harder question is whether you can *run* one. A 1B-parameter model needs 4 GB just for fp32 weights, and that’s before activations, gradients, or the KV cache. Every frontier deployment is a fight against memory, latency, and energy. **Milestone 06** (MLPerf) is the optimization arc — profiling, compression, and acceleration — that turns the architecture you just validated into something that actually ships.



---

PART V

# Optimization Tier

---



## 🔥 Chapter 24

# Module 14: Profiling

You cannot optimize what you have not measured. Profiling is the systems skill that turns “I think this is slow” into “this layer spends 62% of its time waiting on HBM reads at 5% of peak GFLOP/s.” Before you reach for quantization, kernel fusion, or KV caching, you need numbers: parameter count, FLOPs per forward pass, peak activation memory, median latency, and an arithmetic-intensity reading on the roofline. This module builds those instruments end-to-end so every subsequent optimization module has ground truth to point at.

**i** Module Info

**OPTIMIZATION TIER** | Difficulty: ●●○○ | Time: 3-5 hours | Prerequisites: 01-13

**Prerequisites:** Modules 01-13 means you should have:

- Built the complete ML stack (Modules 01-08)
- Implemented CNN architectures (Module 09) or Transformers (Modules 10-13)
- Models to profile and optimize

**Why these prerequisites:** You’ll profile models built in Modules 1-13. Understanding the implementations helps you interpret profiling results — for example, why attention is memory-bound.

## 24.1 Overview

You have built a working ML framework. Now you have to make it fast. The Optimization Tier starts here, and it starts with a rule that almost every engineer breaks at least once: **measure before you optimize**. Guess at the bottleneck and you will spend a week speeding up code that was never on the critical path.

This module gives you the instruments. You’ll build a profiler that counts parameters, estimates FLOPs, tracks memory, and measures latency with enough statistical rigor that the numbers actually mean something. By the end you can answer the questions every optimization decision rests on: Is this model compute-bound or memory-bound? Which layer dominates? Where will quantization or caching pay off — and where will it waste your time?

Every later module in this tier — quantization, compression, acceleration, KV-caching — depends on the data this profiler produces. Build the instrument first. Then optimize.

## 24.2 The Optimization Tier Flow

Profiling (Module 14) is the gateway to the Optimization tier, which follows **Measure** → **Transform** → **Validate**:

Profiling (14) → Model-Level (15-16) → Runtime (17-18) → Benchmarking (19)

↓

↓

↓

↓

"What's slow?"

"Shrink the model"

"Speed up execution"

"Did it work?"

**Model-Level Optimizations (15-16)** — change the model itself:

- Quantization: FP32 → INT8 for 4× compression

- Compression: Prune unnecessary weights

**Runtime Optimizations (17-18)** — change how execution happens:

- Acceleration: Vectorization, kernel fusion (general-purpose)
- Memoization: KV-cache for transformers (domain-specific)

You can't optimize what you can't measure. Profiling comes first because every other tier depends on its output.

## 24.3 Learning Objectives

💡 By completing this module, you will:

- **Implement** a comprehensive Profiler class that measures parameters, FLOPs, memory, and latency
- **Analyze** performance characteristics to identify compute-bound vs memory-bound workloads
- **Master** statistical measurement techniques with warmup runs and outlier handling
- **Connect** profiling insights to optimization opportunities in quantization, compression, and caching

## 24.4 What You'll Build




Figure 24.1: **TinyTorch Profiling System:** Tools for measuring execution time and memory allocation.

### Implementation roadmap:

Table 24.1 lays out the implementation in order, one part at a time.

Table 24.1: **Implementation roadmap for the Profiler class and its measurement methods.**

| Step | What You'll Implement           | Key Concept                             |
|------|---------------------------------|-----------------------------------------|
| 1    | <code>count_parameters()</code> | Model size and memory footprint         |
| 2    | <code>count_flops()</code>      | Computational cost estimation           |
| 3    | <code>measure_memory()</code>   | Activation and gradient memory tracking |
| 4    | <code>measure_latency()</code>  | Statistical timing with warmup          |



| Step | What You'll Implement                | Key Concept                        |
|------|--------------------------------------|------------------------------------|
| 5    | <code>profile_forward_pass()</code>  | Comprehensive performance analysis |
| 6    | <code>profile_backward_pass()</code> | Training cost estimation           |

The pattern you'll enable:

```
# Comprehensive model analysis for optimization decisions
profiler = Profiler()
profile = profiler.profile_forward_pass(model, input_data)
print(f"Bottleneck: {profile['bottleneck']}") # "memory" or "compute"
```

### 24.4.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- GPU profiling (we measure CPU performance with NumPy)
- Distributed profiling (that's for multi-GPU setups)
- CUDA kernel profilers (PyTorch uses `torch.profiler` for GPU analysis)
- Layer-by-layer visualization dashboards (TensorBoard provides this)

**You are building the measurement foundation.** Visualization and GPU profiling come with production frameworks.

## 24.5 API Reference

This section provides a quick reference for the Profiler class you'll build. Use it while implementing and debugging.

### 24.5.1 Constructor

```
Profiler()
```

Initializes profiler with measurement tracking structures.

### 24.5.2 Core Methods

Table 24.2 lists the methods in this group.

Table 24.2: Core measurement methods on the Profiler class.

| Method                        | Signature                                                  | Description                                     |
|-------------------------------|------------------------------------------------------------|-------------------------------------------------|
| <code>count_parameters</code> | <code>count_parameters(model) -&gt; int</code>             | Count total trainable parameters                |
| <code>count_flops</code>      | <code>count_flops(model, input_shape) -&gt; int</code>     | Count FLOPs per sample (batch-size independent) |
| <code>measure_memory</code>   | <code>measure_memory(model, input_shape) -&gt; Dict</code> | Measure memory usage components                 |

| Method                       | Signature                                                                         | Description                               |
|------------------------------|-----------------------------------------------------------------------------------|-------------------------------------------|
| <code>measure_latency</code> | <code>measure_latency(model, input_tensor, warmup, iterations) -&gt; float</code> | Measure inference latency in milliseconds |

### 24.5.3 Analysis Methods

Table 24.3 lists the methods in this group.

Table 24.3: **Higher-level analysis methods on the Profiler class.**

| Method                             | Signature                                                          | Description                        |
|------------------------------------|--------------------------------------------------------------------|------------------------------------|
| <code>profile_layer</code>         | <code>profile_layer(layer, input_shape) -&gt; Dict</code>          | Comprehensive single-layer profile |
| <code>profile_forward_pass</code>  | <code>profile_forward_pass(model, input_tensor) -&gt; Dict</code>  | Complete forward pass analysis     |
| <code>profile_backward_pass</code> | <code>profile_backward_pass(model, input_tensor) -&gt; Dict</code> | Training iteration analysis        |

### 24.5.4 Utility Functions

Table 24.4 lists the methods in this group.

Table 24.4: **Utility functions for quick profiling and weight analysis.**

| Function                                 | Signature                                                                 | Description                                        |
|------------------------------------------|---------------------------------------------------------------------------|----------------------------------------------------|
| <code>quick_profile</code>               | <code>quick_profile(model, input_tensor, profiler=None) -&gt; Dict</code> | One-call convenience profiling                     |
| <code>analyze_weight_distribution</code> | <code>analyze_weight_distribution(model, percentiles) -&gt; Dict</code>   | Statistical analysis of model weight distributions |

## 24.6 Core Concepts

This section covers the fundamental ideas you need to understand profiling deeply. Measurement is the foundation of optimization, and understanding what you’re measuring matters as much as how you measure it.

### 24.6.1 Why Profile First

Optimization without measurement is guessing. You spend a week speeding up a layer that wasn’t on the critical path while the real bottleneck sits untouched. Profiling replaces intuition with ground truth: where time and memory actually go, not where you assumed they did.

Take a slow transformer. Is it attention? The feed-forward layers? Matrix multiplication? Memory transfers? Without numbers, you’re picking randomly. With numbers, you might find that 80% of time is in attention and it’s memory-bound — and now you know to reach for FlashAttention rather than a faster matmul kernel.

The workflow is always the same: measure to baseline, analyze to find the bottleneck, optimize the critical path (not every operation), measure again to verify. Repeat until the numbers hit your target. Your profiler implements the measure-and-analyze steps; later modules supply the optimizations.

### 24.6.2 Timing Operations

Accurate timing is harder than it looks in modern systems due to OS variance, cache warmup effects, and measurement overhead. To counteract these hidden variables, your `measure_latency` method implements a rigorous statistical approach, ensuring the hardware reaches a steady state before any measurements are recorded:

The code in `lst-14-measure-latency` makes this concrete.

```
def measure_latency(self, model, input_tensor, warmup: int = 10, iterations: int =
100) -> float:
    """Measure model inference latency with statistical rigor."""
    # Warmup runs to stabilize performance
    for _ in range(warmup):
        _ = model.forward(input_tensor)

    # Measurement runs
    times = []
    for _ in range(iterations):
        start_time = time.perf_counter()
        _ = model.forward(input_tensor)
        end_time = time.perf_counter()
        times.append((end_time - start_time) * 1000) # Convert to milliseconds

    # Calculate statistics - use median for robustness
    times = np.array(times)
    median_latency = np.median(times)

    return float(median_latency)
```

: **Listing 14.1 — measure\_latency with warmup and median.** Runs warmup iterations to reach steady state, then times a fixed number of forward passes and returns the median to reject OS-noise outliers. {`lst-14-measure-latency`}

The warmup phase is critical. The first few runs are artificially slow due to cold CPU caches, Python interpreter overhead, and NumPy initialization. Running 10+ warmup iterations forces the system into a steady state, yielding reliable baseline measurements.

Median, not mean. A single OS interrupt or garbage-collection pause during measurement can blow the mean apart; the median ignores it. Median captures typical performance, which is what you want to compare across runs. (For SLA work you also report p95 or p99 — but that’s a separate question.)

### 24.6.3 Memory Profiling

Memory profiling reveals three distinct components: parameter memory (model weights), activation memory (forward pass intermediate values), and gradient memory (backward pass derivatives). Each has different characteristics and optimization strategies.

Here’s how your profiler tracks memory usage:

The code in `lst-14-measure-memory` makes this concrete.

```

def measure_memory(self, model, input_shape: Tuple[int, ...]) -> Dict[str, float]:
    """Measure memory usage during forward pass."""
    # Start memory tracking
    tracemalloc.start()

    # Calculate parameter memory
    param_count = self.count_parameters(model)
    parameter_memory_bytes = param_count * BYTES_PER_FLOAT32
    parameter_memory_mb = parameter_memory_bytes / MB_TO_BYTES

    # Create input and measure activation memory
    dummy_input = Tensor(np.random.randn(*input_shape))
    input_memory_bytes = dummy_input.data.nbytes

    # Estimate activation memory (simplified)
    activation_memory_bytes = input_memory_bytes * 2 # Rough estimate
    activation_memory_mb = activation_memory_bytes / MB_TO_BYTES

    # Run forward pass to measure peak memory usage
    _ = model.forward(dummy_input)

    # Get peak memory
    _current_memory, peak_memory = tracemalloc.get_traced_memory()
    peak_memory_mb = (peak_memory - _baseline_memory) / MB_TO_BYTES

    tracemalloc.stop()

    # Calculate efficiency metrics
    useful_memory = parameter_memory_mb + activation_memory_mb
    memory_efficiency = useful_memory / max(peak_memory_mb, 0.001) # Avoid division
    by zero

    return {
        'parameter_memory_mb': parameter_memory_mb,
        'activation_memory_mb': activation_memory_mb,
        'peak_memory_mb': max(peak_memory_mb, useful_memory),
        'memory_efficiency': min(memory_efficiency, 1.0)
    }

```

: **Listing 14.2 — measure\_memory breakdown.** Uses tracemalloc to capture peak allocation during a forward pass, then separates parameter, activation, and peak footprints. {#lst-14-measure-memory}

Parameter memory is persistent and batch-independent. A 125M-parameter model uses 500 MB (125M × 4 bytes per float32) whether you process one sample or a thousand.

Activation memory scales with batch size. Double the batch, double the activations. This is why training needs far more memory than inference at the same model size.

Gradient memory matches parameter memory exactly. Every parameter has one gradient, so training that 125M model adds another 500 MB on top of the weights — and that’s before the optimizer state.

#### 24.6.4 Bottleneck Identification and The Roofline Model

The single most critical insight a profiler yields is whether a workload is **compute-bound** or **memory-bound**. This classification dictates your entire optimization trajectory. Engineers formalize this relationship using the **Roofline Model**, a visual performance framework that plots a system’s peak compute throughput (the horizontal “roof”) against its memory bandwidth (the sloped “attic”).

A workload’s placement under the roof is determined by its **arithmetic intensity**—the ratio of FLOPs executed per byte of memory accessed. The asymptotic complexity of the operation itself dictates where on the roofline it lands: an element-wise add is  $O(N)$  FLOPs on  $O(N)$  bytes (intensity  $\approx 0.08$  FLOPs/byte, memory-bound), while a dense  $N \times N$  matmul is  $O(N^3)$  FLOPs on  $O(N^2)$  bytes (intensity grows linearly with  $N$ , compute-bound once  $N$  is large). Profiling resolves which regime you’re actually in, not just which regime the asymptotics predict.

**Compute-bound** workloads possess high arithmetic intensity. They reside under the flat roof of the model, limited entirely by the arithmetic logic units (e.g., Tensor Cores or SIMD registers). The hardware has ample data but cannot crunch the numbers fast enough. Optimizations here require dense vectorization, kernel fusion, and lower-precision math (like INT8 or FP8).

**Memory-bound** workloads have low arithmetic intensity, trapped under the sloping memory bandwidth line. The processor’s arithmetic units sit idle, starved of data because the hardware cannot fetch information from High Bandwidth Memory (HBM) fast enough. Embedding lookups (sparse gathers) and autoregressive generation (token-by-token processing) notoriously fall here. Optimizations must ruthlessly target data movement: improving cache locality, exploiting SRAM tiling, and reducing the precision footprint.

Your profiler calculates this exact dynamic: if you register a meager GFLOP/s despite running on hardware with massive theoretical throughput, your arithmetic intensity is too low—you have hit the memory wall.

##### **i** Systems Implication: Reading the Roofline to Pick the Right Optimization

The roofline is not just a diagnostic — it is a decision tree. A workload’s position under the roof tells you which lever to pull and which to leave alone. Sitting on the sloped memory bandwidth line? Reducing FLOPs is wasted effort; the chip is already idle waiting on DRAM. You need kernel fusion (Module 17), quantization to shrink per-byte traffic (Module 15), or caching to eliminate redundant reads (Module 18). Sitting under the flat compute roof? Memory is cheap and arithmetic is scarce; reach for structured sparsity, lower precision, or vectorized kernels that keep the ALUs fed. The single most common performance mistake is optimizing the wrong axis: rewriting a compute-bound matmul for better cache locality buys nothing, and fusing a compute-bound kernel can actively hurt by lengthening the critical path. The profiler’s job is to tell you *which axis* — every subsequent module assumes you have read the roofline correctly.

#### 24.6.5 Profiling Tools

The profiler stands on two Python primitives: `time.perf_counter()` for timing and `tracemalloc` for memory.

`time.perf_counter()` reads the system’s highest-resolution monotonic clock — typically nanosecond precision. It returns wall-clock time, so cache misses, context switches, and every other real-world effect show up in your measurement. That’s a feature, not a bug.

`tracemalloc` tracks every Python allocation with byte-level precision and reports both current and peak usage. Peak is what catches the spike that crashes your run.

Production profilers layer on GPU support (CUDA events, NVTX markers), distributed tracing, and kernel-level analysis. The instruments get fancier; the loop stays the same: measure, analyze, identify the bottleneck, optimize.

## 24.7 Production Context

### 24.7.1 Your Implementation vs. PyTorch

Your TinyTorch Profiler and PyTorch’s profiling tools share the same conceptual foundation. The differences are in implementation detail: PyTorch adds GPU support, kernel-level profiling, and distributed tracing. But the core metrics (parameters, FLOPs, memory, latency) are identical.

Table 24.5 places your implementation side by side with the production reference for direct comparison.

Table 24.5: Feature comparison between TinyTorch Profiler and PyTorch profiling tools.

| Feature                    | Your Implementation              | PyTorch                                   |
|----------------------------|----------------------------------|-------------------------------------------|
| <b>Parameter counting</b>  | Direct tensor size               | <code>model.parameters()</code>           |
| <b>FLOP counting</b>       | Per-layer formulas               | <code>FlopCountAnalysis (fvcore)</code>   |
| <b>Memory tracking</b>     | <code>tracemalloc</code>         | <code>torch.profiler</code> , CUDA events |
| <b>Latency measurement</b> | <code>time.perf_counter()</code> | <code>torch.profiler</code> , NVTX        |
| <b>GPU profiling</b>       | CPU only                         | CUDA kernels, memory                      |
| <b>Distributed</b>         | Single process                   | Multi-GPU, NCCL                           |

### 24.7.2 Code Comparison

The following comparison shows equivalent profiling operations in TinyTorch and PyTorch. Notice how the concepts transfer directly, even though PyTorch provides more sophisticated tooling.

## 24.8 Your TinyTorch

```
from tinytorch.perf.profiling import Profiler

# Create profiler
profiler = Profiler()

# Profile model
params = profiler.count_parameters(model)
flops = profiler.count_flops(model, input_shape)
memory = profiler.measure_memory(model, input_shape)
latency = profiler.measure_latency(model, input_tensor)

# Comprehensive analysis
profile = profiler.profile_forward_pass(model, input_tensor)
print(f"Bottleneck: {profile['bottleneck']}")
print(f"GfLOP/s: {profile['gflops_per_second']:.2f}")
```

## 24.9 PyTorch

```
import torch
from torch.profiler import profile, ProfilerActivity

# Count parameters
params = sum(p.numel() for p in model.parameters())

# Profile with PyTorch profiler
with profile(activities=[ProfilerActivity.CPU, ProfilerActivity.CUDA]) as prof:
    output = model(input_tensor)

# Analyze results
print(prof.key_averages().table(sort_by="cpu_time_total"))

# FLOPs (requires fvcore)
from fvcore.nn import FlopCountAnalysis
flops = FlopCountAnalysis(model, input_tensor)
print(f"FLOPs: {flops.total()}")
```

Let's walk through the comparison:

- **Parameter counting:** Both frameworks count total trainable parameters. TinyTorch uses `count_parameters()`, PyTorch uses `sum(p.numel() for p in model.parameters())`.
- **FLOP counting:** TinyTorch implements per-layer formulas. PyTorch uses the `fvcore` library's `FlopCountAnalysis` for more sophisticated analysis.
- **Memory tracking:** TinyTorch uses `tracemalloc`. PyTorch profiler tracks CUDA memory events for GPU memory analysis.
- **Latency measurement:** TinyTorch uses `time.perf_counter()` with warmup. PyTorch profiler uses CUDA events for precise GPU timing.
- **Analysis output:** Both provide bottleneck identification and throughput metrics. PyTorch adds kernel-level detail and distributed profiling.

### 💡 What's Identical

The profiling workflow: measure parameters, FLOPs, memory, and latency to identify bottlenecks. Production frameworks add GPU support and more sophisticated analysis, but the core measurement principles you're learning here transfer directly.

### 24.9.1 Why Profiling Matters at Scale

The stakes get larger with the model. A few examples from production:

- **GPT-3 (175B parameters)** — 652 GB at FP32. Profiling reveals which layers tolerate INT8, which determines whether the model fits on the deployment hardware at all.
- **BERT training** — 80% of step time in self-attention. That single profiling result is what motivated FlashAttention.
- **Image classification, batch 256** — 12 GB of GPU memory used, of which 10 GB is activations. Profiling points straight at gradient checkpointing.

In each case the profiler did not invent the optimization — it told the engineer which optimization was worth implementing. A single session routinely uncovers  $10\times$  speedups or  $4\times$  memory reductions. The instrument earns its keep on the first run.

## 24.10 Check Your Understanding

### 💡 Check Your Understanding — Profiling

Before moving on, verify you can articulate each of the following:

- How to read a FLOP/byte roofline and identify whether a kernel is compute-bound or memory-bound.
- Why median latency is the honest statistic to report for typical performance, and when p95/p99 matter instead.
- How training memory decomposes into parameters + gradients + optimizer state, and why Adam roughly quadruples the raw model footprint.
- Why warmup iterations are not optional — what JIT, cache, and frequency-scaling effects they absorb before steady-state measurement.

If any of these feels fuzzy, revisit the Core Concepts section (especially the Roofline Model and Timing Operations subsections) before moving on.

Work through the following quantitative exercises. Answers are hidden — try each on paper first.

#### Q1: Parameter Memory Calculation

A transformer model has 12 layers, each with a feed-forward network containing two Linear layers: Linear(768, 3072) and Linear(3072, 768). How much memory do the feed-forward network parameters consume across all layers?

### 💡 Answer

Each feed-forward network:

- First layer:  $(768 \times 3072) + 3072 = 2,362,368$  parameters
- Second layer:  $(3072 \times 768) + 768 = 2,360,064$  parameters
- Total per layer: 4,722,432 parameters

Across 12 layers:  $12 \times 4,722,432 = 56,669,184$  parameters.

Memory:  $56,669,184 \times 4 \text{ bytes} = 226,676,736 \text{ bytes} \approx \mathbf{216 \text{ MB}}$ .

That's just the feed-forward networks. Attention adds more parameters on top.

#### Q2: FLOP Counting and Computational Cost

A Linear(512, 512) layer processes a batch of 64 samples. Your profiler's `count_flops()` method returns FLOPs per sample (batch-size independent). How many FLOPs are required for one sample? For the whole batch, if each sample is processed independently?

### 💡 Answer

Per-sample FLOPs (what `count_flops()` returns):  $512 \times 512 \times 2 = \mathbf{524,288 \text{ FLOPs}}$ .

`count_flops()` is batch-size independent — it returns per-sample FLOPs whether you pass `input_shape=(1, 512)` or `(64, 512)`.

For a batch of 64 samples:  $64 \times 524,288 = 33,554,432$  total FLOPs.

Minimum latency at 50 GFLOP/s:  $33,554,432 \div 50 \times 10^9 = \mathbf{0.67 \text{ ms}}$  for the full batch.



That assumes 100% computational efficiency. Real latency is higher because of memory bandwidth and overhead — which is exactly the kind of gap profiling exposes.

### Q3: Memory Bottleneck Analysis

A model achieves 5 GFLOP/s on hardware with 100 GFLOP/s peak compute. The memory bandwidth is 50 GB/s. Is this workload compute-bound or memory-bound?

#### 💡 Answer

Computational efficiency:  $5 \text{ GFLOP/s} \div 100 \text{ GFLOP/s} = 5\%$  **efficiency**.

That gap is the giveaway: this workload is **memory-bound**. The chip can do 100 GFLOP/s but only manages 5, because most of its time is spent waiting on data transfers.

The right optimization strategy is to cut memory traffic — better cache locality, improved data layout, or kernel fusion. Reducing FLOPs won't help when compute is already idle.

### Q4: Training Memory Estimation

A model has 125M parameters (500 MB at FP32). You're training with the Adam optimizer. What's the total memory requirement during training, including gradients and optimizer state?

#### 💡 Answer

- Parameters: 500 MB
- Gradients: 500 MB (one per parameter)
- Adam momentum: 500 MB (first moment estimates)
- Adam velocity: 500 MB (second moment estimates)

Total:  $500 + 500 + 500 + 500 = 2,000 \text{ MB (2 GB)}$  —  $4\times$  the raw model size, just to train.

That's only the model state. Activations add more memory that scales with batch size, putting a typical training run in the 4–8 GB range. This  $4\times$  factor is why optimizer-state sharding (ZeRO, FSDP) exists.

### Q5: Latency Measurement Statistics

You measure latency 100 times and get: median = 10.5 ms, mean = 12.3 ms, min = 10.1 ms, max = 45.2 ms. Which statistic should you report and why?

#### 💡 Answer

Report the **median (10.5 ms)** as the typical latency.

The mean (12.3 ms) is skewed by the outlier (45.2 ms), likely caused by OS interruption or garbage collection. The median is robust to outliers and represents typical performance.

For production SLA planning, you might also report p95 or p99 latency (95th or 99th percentile) to capture worst-case behavior without being skewed by extreme outliers.

## 24.11 Key Takeaways

- **Measure before you optimize:** intuition about bottlenecks is usually wrong, and a week of profiling saves a month of chasing the wrong layer.
- **Arithmetic intensity decides the axis:** FLOPs-per-byte places a workload on the roofline and tells you whether to attack compute or memory — the two optimizations are almost never both correct.
- **Statistical rigor is not pedantry:** warmup + median + confidence intervals are the difference between a real speedup and a noise-driven illusion.

- **Training memory is dominated by optimizer state, not weights:** Adam multiplies the model footprint  $\sim 4\times$  before activations, which is why optimizer sharding (ZeRO/FSDP) exists.

**Coming next:** Module 15 takes the bottlenecks you just surfaced and attacks the most common one — memory pressure — by compressing FP32 weights to INT8, with your profiler acting as the before/after ground truth.

## 24.12 Further Reading

For students who want to understand the academic foundations and professional practices of ML profiling:

### 24.12.1 Seminal Papers

- **Roofline: An Insightful Visual Performance Model** - Williams et al. (2009). Introduces the roofline model for understanding compute vs memory bottlenecks. Essential framework for performance analysis. [ACM CACM](#)
- **PyTorch Profiler: Performance Analysis Tool** - Ansel et al. (2024). Describes PyTorch’s production profiling infrastructure. Shows how profiling scales to distributed GPU systems. [arXiv](#)
- **MLPerf Inference Benchmark** - Reddi et al. (2020). Industry-standard benchmarking methodology for ML systems. Defines rigorous profiling protocols. [arXiv](#)

### 24.12.2 Additional Resources

- **Tool:** [PyTorch Profiler](#) - Production profiling with GPU support
- **Tool:** [TensorFlow Profiler](#) - Alternative framework’s profiling approach
- **Book:** “Computer Architecture: A Quantitative Approach” - Hennessy & Patterson - Chapter 4 covers memory hierarchy and performance measurement

## 24.13 What’s Next

You now have the instrument. Module 15 picks up the first real optimization it enables: **quantization**.

### Coming Up: Module 15 — Quantization

You’ll cut FP32 weights down to INT8 — a  $4\times$  memory reduction — and use this profiler to answer the question that decides whether quantization is worth applying: *which layers tolerate reduced precision, and which ones break?* Profile first, quantize second, profile again to verify. You’re about to see why this loop is the foundation of every production deployment.

### Preview — how your profiler gets used in future modules:

Table 24.6 traces how this module is reused by later parts of the curriculum.

Table 24.6: How the profiler feeds into optimization-tier modules.

| Module                  | What It Does               | Your Profiler In Action                                         |
|-------------------------|----------------------------|-----------------------------------------------------------------|
| <b>15: Quantization</b> | Reduce precision to INT8   | <code>profile_layer()</code> identifies quantization candidates |
| <b>16: Compression</b>  | Prune and compress weights | <code>count_parameters()</code> measures the compression ratio  |
| <b>17: Acceleration</b> | Vectorize computations     | <code>measure_latency()</code> validates the speedup            |

| Module                  | What It Does           | Your Profiler In Action                                             |
|-------------------------|------------------------|---------------------------------------------------------------------|
| <b>19: Benchmarking</b> | Compare across systems | <code>profile_forward_pass()</code> produces the comparable numbers |

## 24.14 Get Started

### 💡 Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### ⚠️ Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.

`e_latency()` validates speedup |



## 🔥 Chapter 25

# Module 15: Quantization

Quantization does not change the  $O(MNK)$  complexity of a matmul. It shrinks the constant factor: 4x less HBM bandwidth per weight, 4x more values per SIMD register (64 INT8 lanes vs 16 FP32 lanes in AVX-512), and a 4x smaller cache footprint that keeps hot weights on-chip. For the memory-bound workloads that dominate autoregressive decoding and on-device inference, that constant-factor win is what makes a 400 MB model fit on a 512 MB IoT device. This module builds asymmetric INT8 quantization, calibration, and a model-level pass that converts every Linear layer in place.

**i** Module Info

**OPTIMIZATION TIER** | Difficulty: ●●●○ | Time: 4-6 hours | Prerequisites: 01-14

**Prerequisites:** Modules 01-14 means you should have:

- Built the complete foundation (Tensor through Training)
- Implemented profiling tools to measure memory usage
- Understanding of neural network parameters and forward passes
- Familiarity with memory calculations and optimization trade-offs

If you can profile a model's memory usage and explain the cost of FP32 storage, you're ready.

## 25.1 Overview

Models have outgrown the devices that need to run them. BERT-base weighs 420 MB, GPT-2 weighs 5.6 GB, and GPT-3 weighs 652 GB — yet a phone has 4–8 GB of RAM total, shared across every app. Every parameter spends 4 bytes on FP32 precision when 8 bits would suffice. Quantization closes that gap: map FP32 weights to INT8 and a model shrinks 4× with typically less than 1% accuracy loss.

In this module you build the INT8 quantization pipeline end-to-end: the core quantize/dequantize functions, a `QuantizedLinear` layer that wraps a trained `Linear`, calibration that fits scale and zero-point to real activation distributions, and a model-level pass that converts every `Linear` in a `Sequential` in place. By the end, you can take a 400 MB checkpoint and ship a 100 MB version that still works.

The math you implement is the same math TensorFlow Lite, PyTorch Mobile, and ONNX Runtime use to fit models on phones, IoT boards, and edge hardware without ever touching the cloud.

## 25.2 Learning Objectives

**💡** By completing this module, you will:

- **Implement** asymmetric INT8 quantization: scale, zero-point, and the quantize/dequantize round trip for 4× memory reduction.
- **Build** calibration that fits scale and zero-point to a real activation distribution from sample inputs.

- **Reason** about quantization error: where it comes from, how it bounds ( $\pm \text{scale}/2$ ), and why neural networks tolerate it.
- **Connect** your implementation to TensorFlow Lite, PyTorch Mobile, and ONNX Runtime — same math, different kernels.
- **Quantify** the memory–accuracy trade-off across model sizes and quantization choices.

## 25.3 What You'll Build




Figure 25.1: **TinyTorch Quantization System**: Methods for converting models to lower precision.

### Implementation roadmap:

Table 25.1 lays out the implementation in order, one part at a time.

Table 25.1: **Implementation roadmap for INT8 quantization and QuantizedLinear.**

| Step | What You'll Implement          | Key Concept                                       |
|------|--------------------------------|---------------------------------------------------|
| 1    | <code>quantize_int8()</code>   | Scale and zero-point calculation, INT8 mapping    |
| 2    | <code>dequantize_int8()</code> | FP32 restoration with quantization parameters     |
| 3    | <code>QuantizedLinear</code>   | Quantized linear layer with compressed weights    |
| 4    | <code>calibrate()</code>       | Input quantization optimization using sample data |
| 5    | <code>quantize_model()</code>  | Full model conversion and memory comparison       |

### The pattern you'll enable:

```
# Compress a 400MB model to 100MB
quantize_model(model, calibration_data=sample_inputs)
# Now model uses 4X less memory with <1% accuracy loss
```

### 25.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Per-channel quantization (PyTorch supports this for finer-grained precision)
- Mixed precision strategies (keeping sensitive layers in FP16/FP32)
- Quantization-aware training (Module 16: Compression introduces this)
- INT8 GEMM kernels (production uses hardware instructions like AVX-512 VNNI)

**You are building per-tensor asymmetric INT8 quantization.** That is enough to compress a real model 4×; the rest is sharper precision and faster kernels.

## 25.4 API Reference

These are the signatures you have to satisfy. Keep this section open in a side pane while you implement — it's the contract the tests check against.

### 25.4.1 Core Functions

```
quantize_int8(tensor: Tensor) -> Tuple[Tensor, float, int]
```

Convert FP32 tensor to INT8 with calculated scale and zero-point.

```
dequantize_int8(q_tensor: Tensor, scale: float, zero_point: int) -> Tensor
```

Restore INT8 tensor to FP32 using quantization parameters.

### 25.4.2 QuantizedLinear Class

Table 25.2 lists the methods and helpers you will implement.

Table 25.2: Core methods on the `QuantizedLinear` class.

| Method                    | Signature                                           | Description                                   |
|---------------------------|-----------------------------------------------------|-----------------------------------------------|
| <code>__init__</code>     | <code>__init__(linear_layer: Linear)</code>         | Create quantized version of Linear layer      |
| <code>calibrate</code>    | <code>calibrate(sample_inputs: List[Tensor])</code> | Optimize input quantization using sample data |
| <code>forward</code>      | <code>forward(x: Tensor) -&gt; Tensor</code>        | Compute output with quantized weights         |
| <code>memory_usage</code> | <code>memory_usage() -&gt; Dict[str, float]</code>  | Calculate memory savings achieved             |

### 25.4.3 Model Quantization

Table 25.3 lists the methods and helpers you will implement.

Table 25.3: **Model-level quantization helper functions.**

| Function                         | Signature                                                 | Description                                |
|----------------------------------|-----------------------------------------------------------|--------------------------------------------|
| <code>quantize_model</code>      | <code>quantize_model(model, calibration_data=None)</code> | Quantize all Linear layers in-place        |
| <code>analyze_model_sizes</code> | <code>analyze_model_sizes(original, quantized)</code>     | Measure compression ratio and memory saved |

#### 25.4.4 Quantizer Class

```
Quantizer()
```

Object-oriented interface wrapping the standalone quantization functions. Provides a convenient API for milestone scripts and production workflows.

Table 25.4 lists the methods and helpers you will implement.

Table 25.4: **Core methods on the Quantizer convenience class.**

| Method                           | Signature                                                 | Description                               |
|----------------------------------|-----------------------------------------------------------|-------------------------------------------|
| <code>quantize_model</code>      | <code>quantize_model(model, calibration_data=None)</code> | Quantize model via static method          |
| <code>analyze_model_sizes</code> | <code>analyze_model_sizes(original, quantized)</code>     | Compare original vs quantized model sizes |

### 25.5 Core Concepts

Three ideas do all the work in this module: how much precision you actually need (range), how to map FP32 to INT8 without wasting it (scale and zero-point), and how to pick those parameters from real data (calibration). Get these right and the implementation almost writes itself.

#### 25.5.1 Precision and Range

FP32 represents about 4.3 billion distinct values across a range from  $10^{-38}$  to  $10^{38}$ . For inference, that's spectacular overkill: trained weights almost always cluster in a tight band like  $[-3, 3]$ , and the network's accuracy depends on patterns in those weights, not on the 23rd bit of mantissa. Small perturbations get absorbed.

INT8 collapses that continuous range to 256 discrete levels ( $-128$  to  $127$ ). The whole game is *which* 256. A tensor whose values live in  $[-0.5, 0.5]$  should not be quantized with the same step size as one whose values live in  $[-10, 10]$  — the first wastes precision, the second loses everything. Quantization is a per-tensor decision about where to spend resolution.

The storage math is unforgiving. One FP32 parameter is 4 bytes; one INT8 parameter is 1 byte. A 100M-parameter model is the difference between 381 MB (FP32) and 95 MB (INT8). The  $4\times$  ratio is fixed because the bit-width ratio is fixed: 32 down to 8.

Crucially, quantization does *not* change the asymptotic complexity of matrix multiplication — a dense  $M\times N$  by  $N\times K$  matmul is still  $O(MNK)$  FLOPs and  $O(MN + NK + MK)$  bytes of traffic whether operands are FP32 or INT8. What changes is the *constant factor*: INT8 operands use one quarter of the memory bandwidth and pack  $4\times$  more operations into the same SIMD register (64 INT8 lanes vs. 16 FP32 lanes



in AVX-512), yielding a 2–4× wall-clock speedup in practice. For memory-bound workloads like embedding lookups and decode-phase attention, that constant-factor win is the difference between fitting on the device and not.

### 25.5.2 Quantization Schemes

Symmetric quantization uses a linear mapping where FP32 zero maps to INT8 zero (zero-point = 0). This simplifies hardware implementation and works well for weight distributions centered around zero. Asymmetric quantization allows the zero-point to shift, better capturing ranges like [0, 1] or [-1, 3] where the distribution is not symmetric.

Your implementation uses asymmetric quantization for maximum flexibility:

The code in `?lst-15-quantize-int8` makes this concrete.

```
def quantize_int8(tensor: Tensor) -> Tuple[Tensor, float, int]:
    """Quantize FP32 tensor to INT8 using asymmetric quantization."""
    data = tensor.data

    # Step 1: Find dynamic range
    min_val = float(np.min(data))
    max_val = float(np.max(data))

    # Step 2: Handle edge case (constant tensor)
    if abs(max_val - min_val) < EPSILON:
        scale = 1.0
        zero_point = 0
        quantized_data = np.zeros_like(data, dtype=np.int8)
        return Tensor(quantized_data), scale, zero_point

    # Step 3: Calculate scale and zero_point
    scale = (max_val - min_val) / (INT8_RANGE - 1)
    zero_point = int(np.round(INT8_MIN_VALUE - min_val / scale))
    zero_point = int(np.clip(zero_point, INT8_MIN_VALUE, INT8_MAX_VALUE))

    # Step 4: Apply quantization formula
    quantized_data = np.round(data / scale + zero_point)
    quantized_data = np.clip(quantized_data, INT8_MIN_VALUE,
                             INT8_MAX_VALUE).astype(np.int8)

    return Tensor(quantized_data), scale, zero_point
```

: **Listing 15.1 — Asymmetric INT8 quantize function.** Derives scale and zero-point from the tensor’s min/max range, then maps FP32 values into [-128, 127] with rounding and clipping. `{#lst-15-quantize-int8}`

The algorithm finds the minimum and maximum values in the tensor, then calculates a scale that maps this range to [-128, 127]. The zero-point determines which INT8 value represents FP32 zero, ensuring minimal quantization error at zero (important for ReLU activations and sparse patterns).

### 25.5.3 Scale and Zero-Point

The scale parameter determines how large each INT8 step is in FP32 space. A scale of 0.01 means each INT8 increment represents 0.01 in the original FP32 values. Smaller scales provide finer precision but can only represent a narrower range; larger scales cover wider ranges but sacrifice precision.

The zero-point is an integer offset that shifts the quantization range. For a symmetric distribution like  $[-2, 2]$ , the zero-point is 0, mapping FP32 zero to INT8 zero. For an asymmetric range like  $[-1, 3]$ , the zero-point is -64, ensuring the quantization levels are distributed optimally across the actual data range.

Here's how dequantization reverses the process:

```
def dequantize_int8(q_tensor: Tensor, scale: float, zero_point: int) -> Tensor:
    """Dequantize INT8 tensor back to FP32."""
    dequantized_data = (q_tensor.data.astype(np.float32) - zero_point) * scale
    return Tensor(dequantized_data)
```

The formula  $(\text{quantized} - \text{zero\_point}) \times \text{scale}$  inverts the quantization mapping. If you quantized 1.5 to INT8 value 50 with scale 0.02 and zero-point -25, dequantization computes  $(50 - (-25)) \times 0.02 = 1.5$ . The round-trip isn't perfect due to quantization being lossy compression, but the error is bounded by the scale value.

### 25.5.4 Post-Training Quantization

Post-training quantization (PTQ) takes a *trained* FP32 model and quantizes it after the fact — no gradient updates, no extra epochs, no labels required. That's the approach you build here. (The alternative, *quantization-aware training*, simulates quantization noise during the training loop so the model learns to be robust to it; you'll see that in Module 16.) `QuantizedLinear` wraps an existing `Linear` and quantizes its weights immediately, deferring activation quantization until calibration:

The code in `?@lst-15-quantized-linear-init` makes this concrete.

```
class QuantizedLinear:
    """Quantized version of Linear layer using INT8 arithmetic."""

    def __init__(self, linear_layer: Linear):
        """Create quantized version of existing linear layer."""
        self.original_layer = linear_layer

        # Quantize weights
        self.q_weight, self.weight_scale, self.weight_zero_point =
            quantize_int8(linear_layer.weight)

        # Quantize bias if it exists
        if linear_layer.bias is not None:
            self.q_bias, self.bias_scale, self.bias_zero_point =
                quantize_int8(linear_layer.bias)
        else:
            self.q_bias = None
            self.bias_scale = None
            self.bias_zero_point = None

        # Store input quantization parameters (set during calibration)
```

```
self.input_scale = None
self.input_zero_point = None
```

: **Listing 15.2 — QuantizedLinear constructor.** Wraps a trained Linear layer, quantizes weights and biases immediately, and reserves input quantization parameters for calibration. {#lst-15-quantized-linear-init}

During inference, the forward pass dequantizes weights on-the-fly, performs the standard FP32 matrix multiplication, and returns FP32 outputs. While this educational approach clarifies the math, production implementations keep the data in 8-bit format entirely, leveraging specialized INT8 GEMM (general matrix multiply) hardware instructions for maximum speed:

The code in ?@lst-15-quantized-linear-forward makes this concrete.

```
def forward(self, x: Tensor) -> Tensor:
    """Forward pass with quantized computation."""
    # Dequantize weights
    weight_fp32 = dequantize_int8(self.q_weight, self.weight_scale,
                                  self.weight_zero_point)

    # Perform computation (same as original layer)
    result = x.matmul(weight_fp32)

    # Add bias if it exists
    if self.q_bias is not None:
        bias_fp32 = dequantize_int8(self.q_bias, self.bias_scale,
                                     self.bias_zero_point)
        result = Tensor(result.data + bias_fp32.data)

    return result
```

: **Listing 15.3 — QuantizedLinear forward pass.** Dequantizes weights and bias on the fly for the matmul. Production kernels skip this step and run the matmul directly in INT8. {#lst-15-quantized-linear-forward}

### 25.5.5 Calibration Strategy

Weights are easy: their values are fixed, so you can compute scale and zero-point from the tensor itself. Activations are not — their range depends on what data flows through the network. Calibration solves this by running a small batch of representative inputs through the layer and recording the activation distribution, then fitting scale and zero-point to *that*:

The code in ?@lst-15-calibrate makes this concrete.

```
def calibrate(self, sample_inputs: List[Tensor]):
    """Calibrate input quantization parameters using sample data."""
    # Collect all input values
    all_values = []
    for inp in sample_inputs:
        all_values.extend(inp.data.flatten())

    all_values = np.array(all_values)

    # Calculate input quantization parameters
```

```

min_val = float(np.min(all_values))
max_val = float(np.max(all_values))

if abs(max_val - min_val) < EPSILON:
    self.input_scale = 1.0
    self.input_zero_point = 0
else:
    self.input_scale = (max_val - min_val) / (INT8_RANGE - 1)
    self.input_zero_point = int(np.round(INT8_MIN_VALUE - min_val /
        self.input_scale))
    self.input_zero_point = np.clip(self.input_zero_point, INT8_MIN_VALUE,
        INT8_MAX_VALUE)

```

: **Listing 15.4 — Activation calibration.** Aggregates sample inputs to fit input scale and zero-point from an empirical activation range before inference. {#lst-15-calibrate}

Calibration typically requires 100-1000 representative samples. Too few samples might miss important distribution characteristics; too many waste time with diminishing returns. The goal is capturing the typical range of activations the model will see during inference.

## 25.6 Production Context

### 25.6.1 Your Implementation vs. PyTorch

Your quantizer implements the same arithmetic PyTorch ships in production. The differences are at the edges: production supports more *schemes* (per-channel, INT4, mixed precision) and runs on dedicated *kernels* (FBGEMM, QNNPACK) that exploit INT8 hardware instructions you didn't build. The math in the middle is identical.

Table 25.5 places your implementation side by side with the production reference for direct comparison.

Table 25.5: Feature comparison between TinyTorch quantizer and PyTorch Quantization.

| Feature            | Your Implementation          | PyTorch Quantization                       |
|--------------------|------------------------------|--------------------------------------------|
| <b>Algorithm</b>   | Asymmetric INT8 quantization | Multiple schemes (INT8, INT4, FP16, mixed) |
| <b>Calibration</b> | Min/max statistics           | MinMax, histogram, percentile observers    |
| <b>Backend</b>     | NumPy (FP32 compute)         | INT8 GEMM kernels (FBGEMM, QNNPACK)        |
| <b>Speed</b>       | 1× (baseline)                | 2-4× faster with INT8 ops                  |
| <b>Memory</b>      | 4× reduction                 | 4× reduction (same compression)            |
| <b>Granularity</b> | Per-tensor                   | Per-tensor, per-channel, per-group         |

### 25.6.2 Code Comparison

The following comparison shows quantization in TinyTorch versus PyTorch. The APIs are remarkably similar, reflecting the universal nature of the quantization problem.

## 25.7 Your TinyTorch

```
from tinytorch.perf.quantization import quantize_model, QuantizedLinear
from tinytorch.core.layers import Linear, Sequential

# Create model
model = Sequential(
    Linear(784, 128),
    Linear(128, 10)
)

# Quantize to INT8
calibration_data = [sample_batch1, sample_batch2, ...]
quantize_model(model, calibration_data)

# Use quantized model
output = model.forward(x)  # 4X less memory!
```

## 25.8 PyTorch

```
import torch
import torch.quantization as quantization

# Create model
model = torch.nn.Sequential(
    torch.nn.Linear(784, 128),
    torch.nn.Linear(128, 10)
)

# Quantize to INT8
model.qconfig = quantization.get_default_qconfig('fbgemm')
model_prepared = quantization.prepare(model)
# Run calibration
for batch in calibration_data:
    model_prepared(batch)
model_quantized = quantization.convert(model_prepared)

# Use quantized model
output = model_quantized(x)  # 4X less memory!
```

Let's walk through the key differences:

- **Line 1-2 (Import):** TinyTorch uses `quantize_model()` function; PyTorch uses `torch.quantization` module with `prepare/convert` API.
- **Lines 4-7 (Model creation):** Both create identical model architectures. The layer APIs are the same.
- **Lines 9-11 (Quantization):** TinyTorch uses one-step `quantize_model()` with calibration data. PyTorch uses three-step API: `configure (qconfig)`, `prepare` (insert observers), `convert` (replace with quantized ops).

- **Lines 13 (Calibration):** TinyTorch passes calibration data as argument; PyTorch requires explicit calibration loop with forward passes.
- **Lines 15-16 (Inference):** Both use standard forward pass. The quantized weights are transparent to the user.

#### 💡 What's Identical

The core quantization mathematics: scale calculation, zero-point mapping, INT8 range clipping. When you debug PyTorch quantization errors, you'll understand exactly what's happening because you implemented the same algorithms.

### 25.8.1 Why Quantization Matters at Scale

The  $4\times$  number sounds modest until you put it next to the device it has to fit on:

#### i Systems Implication: Escaping the Memory Wall via Reduced Precision

As neural networks scale, they inevitably hit the **memory wall** of the Roofline model. Generating tokens autoregressively is inherently memory-bound because the system must constantly stream gigabytes of weight matrices from HBM to the processing cores for every single token. **Quantization is the ultimate systems hack to bypass this wall.** By compressing FP32 weights down to INT8, we instantly cut the memory footprint and the required memory bandwidth by  $4\times$ . Furthermore, modern architectures leverage specific **SIMD** (Single Instruction, Multiple Data) instructions—like AVX-512 VNNI or specialized Tensor Cores—to process these packed 8-bit integers natively. This means we are not just saving RAM; we are quadrupling our cache capacity, slashing the power consumption of data movement, and forcing a memory-bound workload significantly closer to the compute-bound ceiling.

- **Mobile AI:** Modern smartphones have limited RAM shared across all apps. A quantized BERT (`{python} bert_int8_mb`) fits comfortably; the FP32 version (`{python} bert_mb`) causes severe memory pressure and OS-level cache eviction.
- **Edge computing:** IoT devices often have 512 MB RAM. Quantization enables on-device inference for privacy-sensitive applications (medical devices, security cameras) by fitting massive computational graphs into tiny SRAM footprints.
- **Data centers:** Serving 1000 requests/second requires multiple model replicas. With  $4\times$  memory reduction, you fit  $4\times$  more models per GPU, reducing hardware serving costs by 75% and massively improving aggregate throughput.
- **Battery life:** Moving data is vastly more expensive than computing it. INT8 memory transfers and operations consume fractions of the energy required for FP32 equivalents, extending battery life and reducing thermal throttling.

## 25.9 Check Your Understanding

#### 💡 Check Your Understanding — Quantization

Before moving on, verify you can articulate each of the following:

- Why INT8 inference breaks the memory wall on weight-dominated and embedding-heavy workloads ( $4\times$  bandwidth reduction) even though the matmul complexity class is unchanged.
- How asymmetric quantization picks scale and zero-point from a tensor's min/max, and why that choice bounds the round-trip error by  $\pm \text{scale}/2$ .

- Why activation quantization needs calibration data but weight quantization does not.
- Why theoretical  $4\times$  speedup from  $4\times$  SIMD lanes collapses to  $2\text{--}3\times$  in practice (dequant overhead, memory ceiling, non-GEMM ops).

If any of these feels fuzzy, revisit the Core Concepts section (especially Scale and Zero-Point and Calibration Strategy) before moving on.

Five questions to lock in the trade-offs — memory, precision, calibration, I/O, and hardware. Work them out on paper before unfolding the answers.

### Q1: Memory Calculation

A neural network has three Linear layers:  $784 \rightarrow 256$ ,  $256 \rightarrow 128$ ,  $128 \rightarrow 10$ . How much memory do the weights consume in FP32 vs INT8? Include bias terms.

#### 💡 Answer

##### Parameter count:

- Layer 1:  $(784 \times 256) + 256 = 200,960$
- Layer 2:  $(256 \times 128) + 128 = 32,896$
- Layer 3:  $(128 \times 10) + 10 = 1,290$
- **Total: 235,146 parameters**

##### Memory usage:

- FP32:  $235,146 \times 4 \text{ bytes} = 940,584 \text{ bytes} \approx 0.90 \text{ MB}$
- INT8:  $235,146 \times 1 \text{ byte} = 235,146 \text{ bytes} \approx 0.22 \text{ MB}$
- **Savings: 0.67 MB (75% reduction,  $4\times$  compression)**

The ratio is identical for a model  $1000\times$  larger — that's the point.

### Q2: Quantization Error Bound

For FP32 weights uniformly distributed in  $[-0.5, 0.5]$ , what is the maximum quantization error after INT8 quantization? What is the signal-to-noise ratio in decibels?

#### 💡 Answer

##### Quantization error:

- Range:  $0.5 - (-0.5) = 1.0$
- Scale:  $1.0 / 255 = 0.003922$
- Max error:  $\text{scale} / 2 = \pm 0.001961$  (half a step)

##### Signal-to-noise ratio:

- $\text{SNR} = 20 \times \log_{10}(\text{signal\_range} / \text{quantization\_step})$
- $\text{SNR} = 20 \times \log_{10}(1.0 / 0.003922)$
- $\text{SNR} = 20 \times \log_{10}(255)$
- **$\text{SNR} \approx 48 \text{ dB}$**

Neural networks typically need  $>40 \text{ dB}$ , so INT8 has comfortable headroom. The rule of thumb —  $6 \text{ dB per bit}$  — comes straight out of this calculation: every extra bit doubles the number of levels, and  $20 \cdot \log_{10}(2) \approx 6 \text{ dB}$ .

### Q3: Calibration Strategy

You're quantizing a model for deployment. You have 100,000 calibration samples available. How many should you use, and why? What's the trade-off?

 **Answer**

**Recommended: 100–1000 samples** (typically 500).

**Reasoning:**

- **Too few (<100):** risks missing outliers, producing a scale that clips real activations.
- **Too many (>1000):** diminishing returns; you're recomputing the same min/max.
- **Sweet spot (100–1000):** captures the distribution and finishes in seconds.

**Trade-off analysis:**

- 10 samples: ~1 s, may miss distribution tails → noticeable accuracy drop
- 100 samples: ~5 s, good representation → 98% accuracy
- 1000 samples: ~30 s, comprehensive → 98.5% accuracy
- 10000 samples: ~5 min, overkill → 98.6% accuracy

**Conclusion:** accuracy plateaus around 100–1000 samples. Spend more only when the cost of an error is huge (medical, autonomous vehicles).

#### Q4: Memory Bandwidth Impact

A model has 100M parameters. Loading from SSD to RAM at 500 MB/s, how long does loading take for FP32 vs INT8? How does this affect user experience?

 **Answer**

**Loading time:**

- FP32 size:  $100\text{M} \times 4 \text{ bytes} = 381 \text{ MB}$
- INT8 size:  $100\text{M} \times 1 \text{ byte} = 95 \text{ MB}$
- FP32 load time:  $381 \text{ MB} / 500 \text{ MB/s} = \mathbf{0.8 \text{ seconds}}$
- INT8 load time:  $95 \text{ MB} / 500 \text{ MB/s} = \mathbf{0.19 \text{ seconds}}$
- **Speedup:  $4\times$  faster loading**

**User experience impact:**

- Mobile app launch: 0.8 seconds → 0.19 seconds (**0.6s faster startup**)
- Cloud inference: 0.8 seconds cold-start latency → 0.19 seconds ( **$4\times$  better cold-start throughput**)
- Model updates: 381 MB download → 95 MB download (**75% less data over the wire**)

**Key insight:** the  $4\times$  number isn't just about RAM. It applies to disk reads, network transfers, and cold-start latency — every place a byte has to move.

#### Q5: Hardware Acceleration

Modern CPUs have AVX-512 VNNI instructions that can perform INT8 matrix multiply. How many INT8 operations fit in one 512-bit SIMD register vs FP32? Why might actual speedup be less than this ratio?

 **Answer**

**SIMD capacity:**

- 512-bit register with FP32:  $512 / 32 = \mathbf{16 \text{ values}}$
- 512-bit register with INT8:  $512 / 8 = \mathbf{64 \text{ values}}$
- **Theoretical speedup:  $64 / 16 = 4\times$**

**Why actual speedup is  $2\text{--}3\times$  (not  $4\times$ ):**

1. **Dequantization overhead:** converting INT8 → FP32 for activations costs cycles.



2. **Memory bandwidth ceiling:** INT8 ops are so fast that DRAM can't feed them.
3. **Mixed precision:** activations often stay FP32; only weights are quantized.
4. **Non-GEMM ops:** batch norm, softmax, and friends stay FP32.

**Real-world speedup breakdown:**

- Compute-bound (large matmuls): **3–4× speedup**
- Memory-bound (small layers): **1.5–2× speedup**
- Typical mixed models: **2–3× average speedup**

**Key insight:** INT8 wins biggest when matrix multiplications dominate (transformers, large MLPs). For convolutions with tiny kernels, memory bandwidth caps the gains long before compute does.

## 25.10 Key Takeaways

- **4× bandwidth, not 4× FLOPs:** quantization preserves the  $O(MNK)$  matmul complexity but shrinks every constant (memory, cache footprint, SIMD width) by 4× — that is why it wins on memory-bound workloads.
- **Scale and zero-point are per-tensor decisions:** a single tensor with a poorly fit scale wastes resolution or clips values; calibration makes this data-driven instead of guessed.
- **Weights quantize statically, activations quantize with calibration:** weights are fixed after training; activations depend on input distribution and need 100–1000 representative samples.
- **Hardware lottery applies:** the 4× theoretical speedup becomes 2–3× in practice because dequant overhead, non-GEMM ops, and memory ceilings all tax the gain.

**Coming next:** Module 16 attacks a different axis — instead of shrinking each weight, it removes weights entirely via pruning. Composed with quantization, the stack is ~16× smaller than the FP32 original.

## 25.11 Further Reading

For students who want to understand the academic foundations and production implementations of quantization:

### 25.11.1 Seminal Papers

- **Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference** - Jacob et al. (2018). The foundational paper for symmetric INT8 quantization used in TensorFlow Lite. Introduces quantized training and deployment. [arXiv:1712.05877](#)
- **Mixed Precision Training** - Micikevicius et al. (2018). NVIDIA's approach to training with FP16/FP32 mixed precision, reducing memory and increasing speed. Concepts extend to INT8 quantization. [arXiv:1710.03740](#)
- **Data-Free Quantization Through Weight Equalization and Bias Correction** - Nagel et al. (2019). Techniques for quantizing models without calibration data, using statistical properties of weights. [arXiv:1906.04721](#)
- **ZeroQ: A Novel Zero Shot Quantization Framework** - Cai et al. (2020). Shows how to quantize models without any calibration data by generating synthetic inputs. [arXiv:2001.00281](#)

### 25.11.2 Additional Resources

- **Blog post:** “[Quantization in PyTorch](#)” - Official PyTorch quantization tutorial covering eager mode and FX graph mode quantization

- **Documentation:** [TensorFlow Lite Post-Training Quantization](#) - Production quantization techniques for mobile deployment
- **Survey:** “A Survey of Quantization Methods for Efficient Neural Network Inference” - Gholami et al. (2021) - Comprehensive overview of quantization research. [arXiv:2103.13630](#)

## 25.12 What’s Next

You just made every weight  $4\times$  smaller. The next question is the obvious one: do you need every weight at all? Quantization shrinks values; compression deletes them.

### Coming Up: Module 16 — Compression

Module 16 builds **pruning** — first unstructured (zero out individual weights below a threshold) and then structured (remove whole neurons and channels). Pruning attacks a different axis from quantization: it reduces the *count* of operations, not the *cost* of each one. Composed with what you just built, the two techniques multiply: a pruned-then-quantized model is roughly  $16\times$  smaller than the FP32 original, and noticeably faster too.

#### How quantization composes with what comes next:

Table 25.6 traces how this module is reused by later parts of the curriculum.

Table 25.6: **How quantization stacks with compression, acceleration, and capstone modules.**

| Module                  | What it adds                            | The stack so far                                                                                                               |
|-------------------------|-----------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|
| <b>16: Compression</b>  | Pruning removes redundant weights       | <code>quantize_model(pruned_model)</code> $\rightarrow \sim 16\times$ compression                                              |
| <b>17: Acceleration</b> | Kernel fusion eliminates memory traffic | <code>accelerate(quantized_model)</code> $\rightarrow \sim 8\times$ faster inference                                           |
| <b>20: Capstone</b>     | Deploy the full optimized pipeline      | <code>prune</code> $\rightarrow$ <code>quantize</code> $\rightarrow$ <code>accelerate</code> $\rightarrow$ <code>deploy</code> |

## 25.13 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.

`tinytorch/modules/15_quantization/quantization.ipynb`)\*\* - Run interactively in browser, no setup required - **View Source** - Browse the implementation code :::

 **Save Your Progress**

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## 🔥 Chapter 26

# Module 16: Compression

Compression is the negotiation between model footprint and model quality. Every technique here (pruning, distillation, low-rank approximation) buys memory and bandwidth savings at some cost to accuracy, and hardware only cashes the savings when the sparsity pattern matches what the silicon can skip. The ratio is the thing; this module makes it measurable.

### Module Info

**OPTIMIZATION TIER** | Difficulty: ●●●○ | Time: 5-7 hours | Prerequisites: 01-14

**Prerequisites:** **Modules 01-14** means you should have:

- Built tensors, layers, and the complete training pipeline (Modules 01-08)
- Implemented profiling tools to measure model characteristics (Module 14)
- Comfort with weight distributions, parameter counting, and memory analysis

If you can profile a model's parameters and understand weight distributions, you're ready.

## 26.1 Overview

A modern language model takes 100GB of storage. A phone gives you less than 1GB. A microcontroller gives you under 100KB. The model that wins the benchmark and the model that ships are not the same artifact — and the gap between them is what compression closes.

The trick is that trained networks are mostly dead weight. Studies routinely find that 80–90% of a model's parameters can be removed or approximated with negligible loss in accuracy. In this module you implement four techniques that exploit that redundancy: magnitude pruning zeroes the smallest weights, structured pruning drops entire channels so the hardware can actually skip them, knowledge distillation trains a small student to mimic a large teacher, and low-rank approximation factors weight matrices through SVD. By the end you can take a working model and produce a smaller one, while reasoning about where the accuracy went.

## 26.2 Learning Objectives

### By completing this module, you will:

- **Implement** magnitude-based pruning to remove 80-90% of small weights while preserving accuracy
- **Master** structured pruning that creates hardware-friendly sparsity patterns by removing entire channels
- **Build** knowledge distillation systems that compress models 10x through teacher-student training

- **Understand** compression trade-offs between sparsity ratio, inference speed, memory footprint, and accuracy preservation
- **Analyze** when to apply different compression techniques based on deployment constraints and performance requirements

## 26.3 What You'll Build




Figure 26.1: **TinyTorch Compression System:** Reducing model size via pruning and distillation.

### Implementation roadmap:

Table 26.1 lays out the implementation in order, one part at a time.

Table 26.1: **Implementation roadmap for pruning and knowledge distillation.**

| Step | What You'll Implement               | Key Concept                          |
|------|-------------------------------------|--------------------------------------|
| 1    | <code>measure_sparsity()</code>     | Calculate percentage of zero weights |
| 2    | <code>magnitude_prune()</code>      | Remove weights below threshold       |
| 3    | <code>structured_prune()</code>     | Remove entire channels by importance |
| 4    | <code>KnowledgeDistillation</code>  | Train small model from large teacher |
| 5    | <code>low_rank_approximate()</code> | Compress matrices via SVD            |

### The pattern you'll enable:

```
# Compress a model by removing 80% of smallest weights
magnitude_prune(model, sparsity=0.8)
sparsity = measure_sparsity(model) # Returns ~80%
```

### 26.3.1 What You're NOT Building

To keep the module focused, you will **not** implement:

- Sparse storage formats like CSR (production frameworks delegate to `scipy.sparse`)
- Iterative prune-and-fine-tune schedules

- Dynamic pruning during training via hooks and callbacks
- Joint quantization + pruning pipelines

You are building the **algorithms** that decide which weights to remove and how. Squeezing real wall-clock speedup out of the resulting sparsity is the job of the kernels in the next module.

## 26.4 API Reference

This section provides a quick reference for the compression functions and classes you'll build. Use it as your guide while implementing and debugging.

### 26.4.1 Sparsity Measurement

```
measure_sparsity(model) -> float
```

Calculate the percentage of zero weights in a model. Essential for tracking compression effectiveness.

### 26.4.2 Pruning Methods

Table 26.2 lists the functions you will implement.

Table 26.2: **Magnitude and structured pruning functions.**

| Function                      | Signature                                             | Description                                        |
|-------------------------------|-------------------------------------------------------|----------------------------------------------------|
| <code>magnitude_prune</code>  | <code>magnitude_prune(model, sparsity=0.9)</code>     | Remove smallest weights to achieve target sparsity |
| <code>structured_prune</code> | <code>structured_prune(model, prune_ratio=0.5)</code> | Remove entire channels based on L2 norm importance |

### 26.4.3 Knowledge Distillation

```
KnowledgeDistillation(teacher_model, student_model, temperature=3.0, alpha=0.7)
```

**Constructor Parameters:** - `teacher_model`: Large pre-trained model with high accuracy - `student_model`: Smaller model to train via distillation - `temperature`: Softening parameter for probability distributions (typical: 3-5) - `alpha`: Weight for soft targets (0.7 = 70% teacher, 30% hard labels)

**Key Methods:**

Table 26.3 lists the functions you will implement.

Table 26.3: **Key method on the knowledge-distillation helper class.**

| Method                         | Signature                                                                               | Description                        |
|--------------------------------|-----------------------------------------------------------------------------------------|------------------------------------|
| <code>distillation_loss</code> | <code>distillation_loss(student_logits, teacher_logits, true_labels)</code><br>-> float | Combined soft and hard target loss |

### 26.4.4 Low-Rank Approximation

```
low_rank_approximate(weight_matrix, rank_ratio=0.5) -> Tuple[ndarray, ndarray,
ndarray]
```

**Parameters:** - `weight_matrix`: Weight matrix to compress (e.g., (512, 256) Linear layer weights) - `rank_ratio`: Fraction of original rank to keep (0.5 = keep 50% of singular values)

**Returns:** - `U`: Left singular vectors (shape:  $m \times k$ ) - `S`: Singular values (shape:  $k$ ) - `V`: Right singular vectors (shape:  $k \times n$ )

Where  $k = \text{rank\_ratio} \times \min(m, n)$ . Reconstruct approximation with  $U @ \text{diag}(S) @ V$ .

## 26.5 Core Concepts

This section covers the fundamental ideas you need to understand model compression deeply. These concepts apply across all ML frameworks and deployment scenarios.

### 26.5.1 Pruning Fundamentals

Trained networks are over-parameterized: 50–90% of their weights contribute almost nothing to predictions. Pruning exploits that slack by zeroing the deadweight, leaving a sparse network that computes nearly the same function with a fraction of the parameters.

The signal you key off is magnitude. When you compute  $y = W @ x$ , a weight of 0.001 barely moves the output compared to one of magnitude 2.0 or 3.0. Find the small ones, zero them, and most of the network's behavior survives.

Here's how your magnitude pruning implementation works:

The code in `?@lst-16-compression-magnitude-prune` makes this concrete.

```
def magnitude_prune(model, sparsity=0.9):
    """Remove weights with smallest magnitudes to achieve target sparsity."""
    # Collect all weights from model (excluding biases)
    all_weights = []
    weight_params = []

    for param in model.parameters():
        if len(param.shape) > 1: # Only weight matrices, not bias vectors
            all_weights.extend(param.data.flatten())
            weight_params.append(param)

    # Calculate magnitude threshold at desired percentile
    magnitudes = np.abs(all_weights)
    threshold = np.percentile(magnitudes, sparsity * 100)

    # Apply pruning mask: zero out weights below threshold
    for param in weight_params:
        mask = np.abs(param.data) >= threshold
        param.data = param.data * mask # In-place zeroing

    return model
```



: Listing 16.1 — Magnitude pruning via a global percentile threshold. {#lst-16-compression-magnitude-prune}

The elegance is in the percentile-based threshold. Setting `sparsity=0.9` means “remove the bottom 90% of weights by magnitude.” NumPy’s `percentile` function finds the exact value that splits the distribution, and then a binary mask zeros out everything below that threshold.

To understand why this works, consider a typical weight distribution after training:

Weight Magnitudes (sorted):

```
[0.001, 0.002, 0.003, ..., 0.085, 0.087, ..., 0.95, 1.2, 2.3, 3.1]
```

90%

10%

Small, removable

Large, important

```
90th percentile = 0.087
```

```
Threshold mask: magnitude >= 0.087
```

```
Result: Keep only weights >= 0.087 (top 10%)
```

The critical insight is that weight distributions in trained networks are heavily skewed toward zero. Most weights contribute minimally, so removing them preserves the essential computation while dramatically reducing storage and compute.

The memory impact is immediate. A model with 10 million parameters at 90% sparsity has only 1 million active weights. With sparse storage formats (like scipy’s CSR matrix), that translates to a 90% memory reduction. The compute savings come from skipping zero multiplications, but caching them in requires sparse-computation libraries — which is exactly why structured pruning exists.

**Complexity of pruning itself:** finding the global magnitude threshold is  $O(N \log N)$  for a sort (or  $O(N)$  with `quickselect / np.partition`) across  $N$  total weights — a one-shot cost that dominates nothing. The *compression ratio* as a function of sparsity level  $s$  (fraction removed) is simply  $1/(1-s)$ : at  $s=0.5$  you halve the storage, at  $s=0.9$  you get  $10\times$  compression, at  $s=0.99$  you get  $100\times$ . But the *compute* ratio only tracks this if the sparse kernel can actually skip the zeros — and on dense accelerators it cannot, which motivates the structured/unstructured distinction below.

## 26.5.2 Structured vs Unstructured Pruning

Magnitude pruning creates unstructured sparsity: zeros scattered randomly throughout weight matrices. This achieves high compression ratios but creates irregular memory access patterns that modern hardware struggles to accelerate. Structured pruning solves this by removing entire computational units like channels, neurons, or attention heads.

Think of the difference like editing text. Unstructured pruning removes random letters from words, making them hard to read quickly. Structured pruning removes entire words or sentences, preserving readability while reducing length.

Unstructured Sparsity (Magnitude Pruning):

```
Channel 0: [2.1, 0.0, 1.8, 0.0, 3.2]    ← Scattered zeros
Channel 1: [0.0, 2.8, 0.0, 2.1, 0.0]    ← Irregular pattern
Channel 2: [1.5, 0.0, 2.4, 0.0, 1.9]    ← Hard to optimize
```

Structured Sparsity (Channel Pruning):

```
Channel 0: [2.1, 1.3, 1.8, 0.9, 3.2]    ← Fully dense
Channel 1: [0.0, 0.0, 0.0, 0.0, 0.0]    ← Fully removed
Channel 2: [1.5, 2.2, 2.4, 1.1, 1.9]    ← Fully dense
```

Structured pruning requires deciding which channels to remove. Your implementation uses L2 norm as an importance metric:

The code in `?lst-16-compression-structured-prune` makes this concrete.

```
def structured_prune(model, prune_ratio=0.5):
    """Remove entire channels based on L2 norm importance."""
    for layer in model.layers:
        if isinstance(layer, Linear):
            weight = layer.weight.data

            # Calculate L2 norm for each output channel (column)
            channel_norms = np.linalg.norm(weight, axis=0)

            # Find channels to prune (lowest importance)
            num_channels = weight.shape[1]
            num_to_prune = int(num_channels * prune_ratio)

            if num_to_prune > 0:
                # Get indices of smallest channels
                prune_indices = np.argpartition(channel_norms,
                                                num_to_prune)[:num_to_prune]

                # Zero out entire channels
                weight[:, prune_indices] = 0

                # Also zero corresponding bias elements
                if layer.bias is not None:
                    layer.bias.data[prune_indices] = 0

    return model
```

: Listing 16.2 — Structured pruning by L2 norm of output channels. `{#lst-16-compression-structured-prune}`

The L2 norm  $\|W[:, i]\|_2 = \sqrt{\sum(w_{ij}^2)}$  measures the total magnitude of a channel. Channels with small norms contribute less to the output, so dropping them costs the least. Removing whole channels also creates *block* sparsity, which hardware can vectorize over the channels that remain.

The key insight in structured pruning is that you remove entire computational units, not scattered weights. When you zero out channel  $i$  in layer  $l$ , you’re eliminating: - All connections from that channel to the next layer (forward propagation) - All gradient computation for that channel (backward propagation) - The entire channel’s activation storage (memory footprint reduction)

At first glance, structured pruning might seem inferior because it achieves lower overall compression than magnitude pruning. However, the theoretical parameter reduction of unstructured sparsity rarely translates to actual inference speedups on modern hardware.

### **i** Systems Implication: The Hardware Lottery

Why accept lower sparsity ratios with structured pruning? It comes down to the “Hardware Lottery.” Modern AI accelerators—whether NVIDIA GPUs, Google TPUs, or edge NPUs—are fundamentally designed for dense matrix multiplication. They rely on contiguous memory accesses and highly par-

allel vector engines. These architectures cannot efficiently skip scattered zeros without incurring complex, slow indexing overhead that destroys performance. Structured pruning, by removing entire channels, creates smaller but perfectly dense matrices. Instead of forcing the hardware to adapt to an irregular algorithm, you are adapting the algorithm to win the hardware lottery, ensuring the underlying silicon executes the remaining work at peak efficiency.

Because structured pruning preserves dense computation, it directly enables: 1. **Memory Coalescing:** Hardware can fetch the remaining dense channels sequentially, maximizing memory bandwidth utilization. 2. **SIMD Vectorization:** CPUs and GPUs can process multiple channels in perfect parallel lockstep. 3. **Zero Indexing Overhead:** Execution does not require specialized sparse matrix formats (like CSR) to track non-zero locations. 4. **Cache Spatial Locality:** Sequential memory access perfectly aligns with L1/L2 cache line prefetching.

The ultimate trade-off is clear: structured pruning achieves lower sparsity (typically 30-50%) than magnitude pruning (80-90%), but the block sparsity it yields translates directly to proportional hardware acceleration. On dense accelerators, structured sparsity routinely provides 2-3x end-to-end speedups, whereas unstructured sparsity often results in zero acceleration—or even slowdowns—without bespoke sparse computational kernels.

### 26.5.3 Knowledge Distillation

Pruning and SVD shrink an existing model in place. Distillation does something stranger: it builds a brand-new, smaller model and trains it to imitate the bigger one. The large “teacher” never ships — it only exists to supervise the compact “student” that does.

The trick is the *kind* of supervision. Standard training uses one-hot labels — for a cat image, [0, 0, 1, 0] (100% cat). The teacher’s predictions are softer: [0.02, 0.05, 0.85, 0.08] (85% cat, with some mass on visually similar classes). Those off-diagonal probabilities are the teacher’s tacit knowledge about class similarity, and they teach the student more per gradient step than a hard label ever could.

Temperature scaling controls how soft the distributions become:

The code in `?@lst-16-compression-distillation-loss` makes this concrete.

```
def _softmax(self, logits):
    """Compute softmax with numerical stability."""
    exp_logits = np.exp(logits - np.max(logits, axis=-1, keepdims=True))
    return exp_logits / np.sum(exp_logits, axis=-1, keepdims=True)

def distillation_loss(self, student_logits, teacher_logits, true_labels):
    """Calculate combined distillation loss."""
    # Soften distributions with temperature
    student_soft = self._softmax(student_logits / self.temperature)
    teacher_soft = self._softmax(teacher_logits / self.temperature)

    # Soft target loss (KL divergence)
    soft_loss = self._kl_divergence(student_soft, teacher_soft)

    # Hard target loss (cross-entropy)
    student_hard = self._softmax(student_logits)
    hard_loss = self._cross_entropy(student_hard, true_labels)

    # Combined loss
```

```
total_loss = self.alpha * soft_loss + (1 - self.alpha) * hard_loss

return total_loss
```

: Listing 16.3 — Temperature-scaled distillation loss combining soft and hard targets. {#lst-16-compression-distillation-loss}

Dividing logits by temperature before softmax spreads probability mass across classes. At `temperature=1` you get the usual peaked softmax. At `temperature=3` the distribution flattens, exposing the teacher’s relative uncertainty between classes. The student learns both what the teacher picks *and* what it almost picked.

The combined loss balances two pressures. The soft loss (weighted `alpha=0.7`) pushes the student to imitate the teacher’s reasoning. The hard loss (weighted `1-alpha=0.3`) keeps it anchored to the ground truth. With this recipe, students routinely match teachers at 10x compression with 2–5% accuracy loss.

### 26.5.4 Low-Rank Approximation Theory

Weight matrices in trained networks are usually closer to low-rank than they look. Singular Value Decomposition (SVD) gives you the mathematically optimal way to approximate any matrix with fewer parameters, ranked by how much variance each direction explains.

The mechanism is matrix factorization. Instead of storing a full (512, 256) weight matrix with 131,072 parameters, you decompose it into two thinner matrices that capture the essential structure:

The code in `?@lst-16-compression-low-rank` makes this concrete.

```
def low_rank_approximate(weight_matrix, rank_ratio=0.5):
    """Approximate weight matrix using SVD-based low-rank decomposition."""
    m, n = weight_matrix.shape

    # Perform SVD: W = U @ diag(S) @ V
    U, S, V = np.linalg.svd(weight_matrix, full_matrices=False)

    # Keep only top-k singular values
    max_rank = min(m, n)
    target_rank = max(1, int(rank_ratio * max_rank))

    # Truncate to target rank
    U_truncated = U[:, :target_rank]      # (m, k)
    S_truncated = S[:target_rank]         # (k, )
    V_truncated = V[:target_rank, :]     # (k, n)

    return U_truncated, S_truncated, V_truncated
```

: Listing 16.4 — Low-rank approximation via truncated SVD. {#lst-16-compression-low-rank}

SVD identifies the principal “directions” in the weight matrix through singular values. Larger singular values capture more variance, so keeping the top *k* preserves most of the matrix’s information at a fraction of the parameter cost.

For a (512, 256) matrix with `rank_ratio=0.5`:

- Original:  $512 \times 256 = 131,072$  parameters
- Compressed:  $(512 \times 128) + 128 + (128 \times 256) = 98,432$  parameters
- Compression ratio: 1.33x (25% reduction)

The win compounds with larger matrices. For a (1024, 1024) matrix at `rank_ratio=0.1`:

- Original: 1,048,576 parameters
- Compressed:  $(1024 \times 102) + 102 + (102 \times 1024) = 208,998$  parameters
- Compression ratio: 5.0x (80% reduction)

The pattern: SVD pays off most where the matrix is wide and the intrinsic rank is low. The reconstruction error you absorb is exactly the sum of squares of the singular values you threw away — pick `rank_ratio` by looking at the singular value spectrum, not by guessing.

### 26.5.5 Compression Trade-offs

Every compression technique trades accuracy for efficiency, but different techniques make different trade-offs. Understanding these helps you choose the right approach for your deployment constraints.

Table 26.4 compares compression, accuracy, and speedup trade-offs across techniques.

Table 26.4: **Compression, accuracy, and speedup trade-offs across compression techniques.**

| Technique                     | Compression Ratio | Accuracy Loss | Hardware Speedup            | Training Required      |
|-------------------------------|-------------------|---------------|-----------------------------|------------------------|
| <b>Magnitude Pruning</b>      | 5-10x             | 1-3%          | Minimal (needs sparse libs) | No (prune pretrained)  |
| <b>Structured Pruning</b>     | 2-3x              | 2-5%          | 2-3x (hardware-friendly)    | No (prune pretrained)  |
| <b>Knowledge Distillation</b> | 10-50x            | 5-10%         | Proportional to size        | Yes (train student)    |
| <b>Low-Rank Approximation</b> | 2-5x              | 3-7%          | Minimal (depends on impl)   | No (SVD decomposition) |

The systems insight is that compression ratio alone does not determine deployment success. A 10x compressed model from magnitude pruning can run *slower* than a 3x compressed model from structured pruning, because the hardware cannot accelerate scattered zeros. Distillation buys the best size-for-accuracy ratio, but only if you have the training budget to retrain a student. Pick the technique that matches the bottleneck — memory, latency, or training time — not the one with the prettiest sparsity number.

## 26.6 Production Context

### 26.6.1 Your Implementation vs. PyTorch

Your TinyTorch compression functions and PyTorch’s pruning utilities share the same core algorithms. The differences are in integration depth: PyTorch provides hooks for pruning during training, automatic mask management, and integration with quantization-aware training. But the fundamental magnitude-based and structured pruning logic is identical.

Table 26.5 places your implementation side by side with the production reference for direct comparison.

Table 26.5: **Feature comparison between TinyTorch compression and PyTorch pruning utilities.**

| Feature                  | Your Implementation             | PyTorch                                        |
|--------------------------|---------------------------------|------------------------------------------------|
| <b>Magnitude Pruning</b> | Global threshold via percentile | <code>torch.nn.utils.prune.l1_unstructu</code> |

| Feature                       | Your Implementation     | PyTorch                                         |
|-------------------------------|-------------------------|-------------------------------------------------|
| <b>Structured Pruning</b>     | L2 norm channel removal | <code>torch.nn.utils.prune.l1_structured</code> |
| <b>Knowledge Distillation</b> | Manual loss calculation | User-implemented (same approach)                |
| <b>Sparse Execution</b>       | Dense NumPy arrays      | Sparse tensors + kernels                        |
| <b>Pruning Schedules</b>      | One-shot pruning        | Iterative + fine-tuning                         |

### 26.6.2 Code Comparison

The following comparison shows equivalent compression operations in TinyTorch and PyTorch. Notice how the core concepts translate directly while PyTorch provides additional automation for production workflows.

## 26.7 Your TinyTorch

```
from tinytorch.perf.compression import magnitude_prune, measure_sparsity

# Create model
model = Sequential(Linear(100, 50), ReLU(), Linear(50, 10))

# Apply magnitude pruning
magnitude_prune(model, sparsity=0.8)

# Measure results
sparsity = measure_sparsity(model) # Returns 80.0 (percentage)
print(f"Sparsity: {sparsity:.1f}%")
```

## 26.8 PyTorch

```
import torch
import torch.nn as nn
import torch.nn.utils.prune as prune

# Create model
model = nn.Sequential(
    nn.Linear(100, 50),
    nn.ReLU(),
    nn.Linear(50, 10)
)

# Apply magnitude pruning
for module in model.modules():
    if isinstance(module, nn.Linear):
        prune.l1_unstructured(module, name='weight', amount=0.8)
        prune.remove(module, 'weight') # Make pruning permanent
```

```
# Measure results
total_params = sum(p.numel() for p in model.parameters())
zero_params = sum((p == 0).sum().item() for p in model.parameters())
sparsity = zero_params / total_params
print(f"Sparsity: {sparsity:.1%}")
```

Let's walk through the key differences:

- **Line 1 (Import):** TinyTorch provides compression in a dedicated `perf.compression` module. PyTorch's `torch.nn.utils.prune` offers similar functionality with additional hooks.
- **Line 4-5 (Model):** Both create identical model architectures. PyTorch's `nn.Sequential` matches TinyTorch's explicit layer composition.
- **Line 8 (Pruning):** TinyTorch uses a simple function call that operates on the entire model. PyTorch requires iterating over modules and applying pruning individually, offering finer control.
- **Line 13 (Permanence):** TinyTorch immediately zeros weights. PyTorch uses masks that can be removed or made permanent, enabling experimentation with different sparsity levels.
- **Line 16-19 (Measurement):** TinyTorch provides a dedicated `measure_sparsity()` function. PyTorch requires manual counting, giving you full control over what counts as "sparse."

#### 💡 What's Identical

The core algorithms for magnitude thresholding, L2 norm channel ranking, and knowledge distillation loss are identical. When you understand TinyTorch compression, you understand PyTorch compression. The production differences are in automation, not algorithms.

### 26.8.1 Why Compression Matters at Scale

The deployment constraints make the case for compression more bluntly than any benchmark:

- **Mobile apps** — models must fit under ~10MB for download and ~50MB for runtime memory
- **Edge devices** — a Raspberry Pi 4 has 4GB total RAM, shared with the OS and every other process
- **Cloud cost** — GPT-class inference costs millions per month at scale; 10x compression is 10x cheaper
- **Latency targets** — self-driving stacks need <100ms end-to-end; compression buys most of that budget
- **Energy** — phone batteries are ~3000mAh, and bigger models drain them faster

A 100 MB model pruned to 90% sparsity becomes 10 MB with sparse storage, which fits the mobile budget. Distilled down to a 1 MB student, it also runs 10x faster, which fits the latency budget. These aren't projections — they're the entry conditions for shipping at all.

## 26.9 Check Your Understanding

### 💡 Check Your Understanding — Compression

Before moving on, verify you can articulate each of the following:

- The Hardware Lottery: why structured sparsity beats unstructured for real-world speedups on dense accelerators, even at lower compression ratios.
- How magnitude pruning picks a threshold from the global weight distribution, and why compression ratio =  $1/(1 - \text{sparsity})$ .

- Why knowledge distillation's soft targets (teacher logits under temperature) teach more than one-hot labels per gradient step.
- When low-rank approximation pays off (wide matrices, low intrinsic rank) and when SVD is wasted effort.

If any of these feels fuzzy, revisit the Core Concepts section (especially Structured vs Unstructured Pruning and Compression Trade-offs) before moving on.

Test yourself with these systems-thinking questions. They're built to sharpen the intuition you'll need every time you compress a model in production.

#### Q1: Sparsity Calculation

A Linear layer with shape (512, 256) undergoes 80% magnitude pruning. How many weights remain active?

##### 💡 Answer

Total parameters:  $512 \times 256 = 131,072$

After 80% pruning: 20% remain active =  $131,072 \times 0.2 = 26,214$  **active weights**

Zeroed weights:  $131,072 \times 0.8 = 104,857$  **zeros**

This is why sparsity creates memory savings — 80% of the parameters are literally zero.

#### Q2: Compression Ratio Analysis

You apply magnitude pruning (90% sparsity) and structured pruning (50% channels) sequentially. What's the final sparsity?

##### 💡 Answer

**Trick question!** Structured pruning zeros entire channels, which may already be partially sparse from magnitude pruning.

Approximation: - After magnitude: 90% sparse  $\rightarrow$  10% active weights - Structured removes 50% of channels  $\rightarrow$  removes 50% of rows/columns - Final active weights  $\approx 10\% \times 50\% = 5\%$  **active**  $\rightarrow$  **95% sparse**

Actual result depends on which channels structured pruning removes. If it removes already-sparse channels, sparsity increases less.

#### Q3: Knowledge Distillation Efficiency

Teacher model: 100M parameters, 95% accuracy, 500ms inference Student model: 10M parameters, 92% accuracy, 50ms inference

What's the compression ratio and speedup?

##### 💡 Answer

**Compression ratio:**  $100\text{M} / 10\text{M} = 10\text{x smaller}$

**Speedup:**  $500\text{ms} / 50\text{ms} = 10\text{x faster}$

**Accuracy loss:**  $95\% - 92\% = 3\%$  **degradation**

Why speedup matches compression: the student has 10x fewer parameters, so roughly 10x fewer ops. Linear scaling.

Is this a good trade? **Yes** — 10x compression for 3% accuracy loss is the sweet spot for mobile deployment.

#### Q4: Low-Rank Decomposition Math



A (1000, 1000) weight matrix gets low-rank approximation with rank=100. Calculate parameter reduction.

#### 💡 Answer

Original:  $1000 \times 1000 = \mathbf{1,000,000}$  parameters

SVD decomposition:  $W \approx U @ \text{diag}(S) @ V$

- U: (1000, 100) = 100,000 parameters
- S: (100,) = 100 parameters (diagonal)
- V: (100, 1000) = 100,000 parameters

Compressed:  $100,000 + 100 + 100,000 = \mathbf{200,100}$  parameters

Compression ratio:  $1,000,000 / 200,100 = \sim \mathbf{5x}$  reduction

Memory savings:  $(1,000,000 - 200,100) \times 4 \text{ bytes} = \mathbf{3.1 MB}$  saved (float32)

#### Q5: Structured vs Unstructured Trade-offs

For mobile deployment with tight latency constraints, would you choose magnitude pruning (90% sparsity) or structured pruning (30% sparsity)? Why?

#### 💡 Answer

**Choose structured pruning (30% sparsity)** despite lower compression.

Reasoning: 1. **Hardware acceleration:** Mobile CPUs/GPUs can execute dense channels 2-3x faster than sparse patterns 2. **Latency guarantee:** Structured sparsity gives predictable speedup; magnitude sparsity needs sparse libraries (often unavailable on mobile) 3. **Real speedup:** 30% structured =  $\sim 1.5x$  actual speedup; 90% magnitude = no speedup without custom kernels 4. **Memory:** Both save memory, but latency requirement dominates

**Production insight:** High sparsity  $\neq$  high speedup. Hardware capabilities matter more than compression ratio for latency-critical applications.

## 26.10 Key Takeaways

- **Compression ratio is  $1/(1 - \text{sparsity})$ :** the arithmetic is trivial; the engineering question is whether the underlying kernel can actually skip zeros.
- **Structured > unstructured on dense hardware:** 30–50% structured sparsity beats 90% unstructured sparsity in wall-clock time, because modern accelerators execute dense tiles and pay indexing overhead for scattered zeros.
- **Distillation compresses by transferring knowledge, not zeroing weights:** the student is a new, smaller network trained on the teacher's soft outputs — worth the extra training budget when size-for-accuracy is the goal.
- **SVD pays off for wide, low-rank matrices:** pick the rank ratio from the singular-value spectrum, not by guessing.

**Coming next:** Module 17 answers the question every pruning pass begs — why isn't the model  $10\times$  faster? — by building the vectorized kernels, cache-aware tiling, and kernel fusion that turn structured sparsity into measurable wall-clock speedup.

## 26.11 Further Reading

For students who want to understand the academic foundations and explore compression techniques further:

### 26.11.1 Seminal Papers

- **Learning both Weights and Connections for Efficient Neural Networks** - Han et al. (2015). Introduced magnitude-based pruning and demonstrated 90% sparsity with minimal accuracy loss. Foundation for modern pruning research. [arXiv:1506.02626](#)
- **The Lottery Ticket Hypothesis** - Frankle & Carlin (2019). Showed that dense networks contain sparse subnetworks trainable to full accuracy from initialization. Changed how we think about pruning and network over-parameterization. [arXiv:1803.03635](#)
- **Distilling the Knowledge in a Neural Network** - Hinton et al. (2015). Introduced knowledge distillation with temperature scaling. Enables training compact models that match large model accuracy. [arXiv:1503.02531](#)
- **Pruning Filters for Efficient ConvNets** - Li et al. (2017). Demonstrated structured pruning by removing entire convolutional filters. Showed that L1-norm ranking identifies unimportant channels effectively. [arXiv:1608.08710](#)

### 26.11.2 Additional Resources

- **Survey:** “Model Compression and Hardware Acceleration for Neural Networks” by Deng et al. (2020) - Comprehensive overview of compression techniques and hardware implications
- **Tutorial:** [PyTorch Pruning Tutorial](#) - See how production frameworks implement these concepts
- **Blog:** “The State of Sparsity in Deep Neural Networks” by Uber Engineering - Practical experiences deploying sparse models at scale

## 26.12 What’s Next

You now have a model that is 10x smaller on paper. Module 17 answers the question that follows immediately: *why isn’t it 10x faster?* Sparsity and low rank only show up in wall-clock time when the underlying kernels know how to skip zeros and reuse memory. Without that, your pruned weights are still being multiplied — by zero, but multiplied.

#### Coming Up: Module 17 — Acceleration

Implement hardware-aware optimizations: vectorized matrix kernels, operator fusion, and cache-friendly tiling. You’ll fuse multiple ops into single passes to cut memory-bandwidth pressure and turn the structured sparsity from this module into measurable speedup.

#### Preview — how your compression gets used in future modules:

Table 26.6 traces how this module is reused by later parts of the curriculum.

Table 26.6: **How compression feeds into acceleration, memoization, and benchmarking.**

| Module                  | What It Does                   | Your Compression In Action                                         |
|-------------------------|--------------------------------|--------------------------------------------------------------------|
| <b>17: Acceleration</b> | Optimize computation kernels   | Structured sparsity enables vectorized operations                  |
| <b>18: Memoization</b>  | Cache repeated computations    | <code>compress_model()</code> before caching for memory efficiency |
| <b>19: Benchmarking</b> | Measure end-to-end performance | Compare dense vs sparse model throughput                           |

## 26.13 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## 🔥 Chapter 27

# Module 17: Acceleration

Most neural networks spend more time waiting on memory than computing. Acceleration is the art of rearranging work so the ALUs never idle: SIMD vectorization, cache tiling, and kernel fusion that keeps intermediates in registers instead of round-tripping through HBM. Here you build the primitives that close the gap between theoretical peak FLOPS and measured throughput on real silicon.

### **i** Module Info

**OPTIMIZATION TIER** | Difficulty: ●●●○ | Time: 5-7 hours | Prerequisites: 01-14

**Prerequisites:** **Modules 01-14** means you need:

- Tensor operations (Module 01) for understanding data structures
- Neural network layers (Module 03) for knowing what to accelerate
- Training loops (Module 08) for understanding the performance context
- Profiling tools (Module 14) for measuring acceleration gains

If you can multiply matrices and understand why matrix multiplication is expensive, you're ready.

## 27.1 Overview

Neural networks spend 90% of their time multiplying matrices. The same workload that trains in three hours on optimized kernels takes a week of naive Python loops — same math, same hardware, two orders of magnitude apart. The gap is paid for entirely in how the code talks to the processor.

This module teaches hardware-aware optimization through vectorization and kernel fusion. You'll exploit SIMD lanes, fix memory access patterns, and eliminate the intermediate arrays that quietly burn most of your bandwidth. By the end you'll know why a naive matmul is 100x slower than an optimized one, and you'll have shipped 2–5x speedups against your own baseline.

Acceleration isn't clever algorithms. It's understanding how processors actually work, then writing code that doesn't fight them.

## 27.2 Why Acceleration Before Memoization?

The Optimization tier splits into **model-level** (15–16) and **runtime** (17–18) work:

- **Model-level** (Quantization, Compression): change the model itself.
- **Runtime** (Acceleration, Memoization): change how execution happens.

Acceleration comes first because it is general — vectorization and kernel fusion apply to every numerical operation a network runs: matmuls, convolutions, attention, activations. Memoization is the opposite: a domain-specific trick that pays off mostly for transformer autoregressive generation. Build the general tools first, then specialize.

## 27.3 Learning Objectives

💡 By completing this module, you will:

- **Implement** vectorized matrix multiplication using optimized BLAS libraries for maximum throughput
- **Master** kernel fusion techniques that eliminate memory bandwidth bottlenecks by combining operations
- **Understand** the roofline model and arithmetic intensity to predict performance bottlenecks
- **Analyze** production acceleration strategies for different deployment scenarios (edge, cloud, GPU)

## 27.4 What You'll Build




Figure 27.1: **TinyTorch Acceleration System:** Vectorized kernels and SIMD optimization.

### Implementation roadmap:

Table 27.1 lays out the implementation in order, one part at a time.

Table 27.1: **Implementation roadmap for vectorized kernels and operator fusion.**

| Part | What You'll Implement            | Key Concept                       |
|------|----------------------------------|-----------------------------------|
| 1    | <code>vectorized_matmul()</code> | SIMD and BLAS optimization        |
| 2    | <code>fused_gelu()</code>        | Memory bandwidth reduction        |
| 3    | <code>unfused_gelu()</code>      | Comparison baseline               |
| 4    | <code>tiled_matmul()</code>      | Cache-aware computation           |
| 5    | Performance analysis             | Roofline and arithmetic intensity |

### The pattern you'll enable:

```
# Fast matrix operations using BLAS
output = vectorized_matmul(x, weights) # 10-100x faster than naive loops
```

```
# Memory-efficient activations
activated = fused_gelu(output) # 60% less memory bandwidth
```

### 27.4.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- GPU kernels (that requires CUDA programming, covered in production frameworks)
- Custom CPU assembly (BLAS libraries already provide this)
- Automatic kernel fusion (compilers like XLA do this automatically)
- Multi-threading control (NumPy handles this via OpenBLAS/MKL)

**You are building the understanding.** Hardware-specific implementations come later.

## 27.5 API Reference

This section provides a quick reference for the acceleration functions you'll build. These functions demonstrate optimization techniques that apply to any ML framework.

### 27.5.1 Vectorized Operations

```
vectorized_matmul(a: Tensor, b: Tensor) -> Tensor
```

High-performance matrix multiplication using optimized BLAS libraries that leverage SIMD instructions and cache blocking.

### 27.5.2 Kernel Fusion

Table 27.2 lists the functions you will implement.

Table 27.2: Fused and baseline GELU kernel functions.

| Function                  | Signature                                         | Description                                          |
|---------------------------|---------------------------------------------------|------------------------------------------------------|
| <code>fused_gelu</code>   | <code>fused_gelu(x: Tensor) -&gt; Tensor</code>   | GELU activation with all operations in single kernel |
| <code>unfused_gelu</code> | <code>unfused_gelu(x: Tensor) -&gt; Tensor</code> | Baseline implementation for comparison               |

### 27.5.3 Cache-Aware Operations

Table 27.3 lists the functions you will implement.

Table 27.3: Cache-aware matrix multiplication function.

| Function                  | Signature                                                                    | Description                           |
|---------------------------|------------------------------------------------------------------------------|---------------------------------------|
| <code>tiled_matmul</code> | <code>tiled_matmul(a: Tensor, b: Tensor, tile_size: int) -&gt; Tensor</code> | Cache-optimized matrix multiplication |

## 27.6 Core Concepts

This section covers the fundamental acceleration techniques that apply to any hardware platform. Understanding these concepts will help you optimize neural networks whether you're targeting CPUs, GPUs, or specialized accelerators.

### 27.6.1 Vectorization with NumPy

Modern processors can execute the same operation on multiple data elements simultaneously through SIMD (Single Instruction, Multiple Data) instructions. A traditional loop processes one element per clock cycle, but SIMD can process 4, 8, or even 16 elements in the same time.

Consider a simple element-wise addition. A naive Python loop visits each element sequentially:

```
# Slow: one element per iteration
for i in range(len(x)):
    result[i] = x[i] + y[i]
```

NumPy's vectorized operations automatically use SIMD when you write `x + y`. The processor loads multiple elements into special vector registers and adds them in parallel. This is why vectorized NumPy code can be 10-100x faster than explicit loops.

Here's how vectorized matrix multiplication works in your implementation:

The code in `?lst-17-acceleration-vectorized-matmul` makes this concrete.

```
def vectorized_matmul(a: Tensor, b: Tensor) -> Tensor:
    """Matrix multiplication using optimized BLAS libraries."""
    # Validate shapes - inner dimensions must match
    if a.shape[-1] != b.shape[-2]:
        raise ValueError(
            f"Matrix multiplication shape mismatch: {a.shape} @ {b.shape}. "
            f"Inner dimensions must match: a.shape[-1]={a.shape[-1]} != "
            f"b.shape[-2]={b.shape[-2]}"
        )

    # NumPy calls BLAS GEMM which uses:
    # - SIMD vectorization for parallel arithmetic
    # - Cache blocking for memory efficiency
    # - Multi-threading on multi-core systems
    result_data = np.matmul(a.data, b.data)

    return Tensor(result_data)
```

: Listing 17.1 — Vectorized matmul delegating to BLAS GEMM. `{#lst-17-acceleration-vectorized-matmul}`

The magic happens inside `np.matmul`. NumPy delegates to BLAS (Basic Linear Algebra Subprograms) libraries like OpenBLAS or Intel MKL. These libraries have been optimized over decades to exploit every hardware feature: SIMD instructions, cache hierarchies, and multiple cores. The same Python code that takes 800ms with naive loops completes in 8ms with BLAS.

### 27.6.2 BLAS and LAPACK

BLAS provides three levels of operations, each with different performance characteristics:



- **Level 1:** Vector operations (AXPY:  $y = \alpha x + y$ ). Memory-bound, low arithmetic intensity.
- **Level 2:** Matrix-vector operations (GEMV:  $y = \alpha Ax + \beta y$ ). Better arithmetic intensity, still memory-limited.
- **Level 3:** Matrix-matrix operations (GEMM:  $C = \alpha AB + \beta C$ ). High arithmetic intensity, compute-bound.

Matrix multiplication (GEMM) dominates neural network training: every linear layer, every attention head, every convolution ultimately reduces to it. GEMM performs  $2N^3$  floating-point operations while reading only  $3N^2$  elements. For a  $1024 \times 1024$  matrix that's 2.1 billion operations on just 12 MB of data — an arithmetic intensity of  $\sim 171$  FLOPs/byte. That ratio is why GEMM is the operation hardware designers tune for first.

### 27.6.3 Memory Layout Optimization

When a processor needs data from main memory, it doesn't fetch individual bytes. It fetches entire cache lines (typically 64 bytes). If your data is laid out sequentially in memory, you get spatial locality: one cache line brings in many useful values. If your data is scattered randomly, every access causes a cache miss and a 100-cycle stall.

Matrix multiplication has interesting memory access patterns. Computing one output element requires reading an entire row from the first matrix and an entire column from the second matrix. Rows are stored sequentially in memory (good), but columns are strided by the matrix width (potentially bad). This is why cache-aware tiling helps:

```
# Cache-aware tiling breaks large matrices into blocks
# Each block fits in cache for maximum reuse
for i_tile in range(0, M, tile_size):
    for j_tile in range(0, N, tile_size):
        for k_tile in range(0, K, tile_size):
            # Multiply tile blocks that fit in L1/L2 cache
            C[i_tile:i_tile+tile_size, j_tile:j_tile+tile_size] +=
                A[i_tile:i_tile+tile_size, k_tile:k_tile+tile_size] @
                B[k_tile:k_tile+tile_size, j_tile:j_tile+tile_size]
```

Your `tiled_matmul` implementation demonstrates this concept, though in practice NumPy's BLAS backend already implements optimal tiling:

The code in `?@lst-17-acceleration-tilled-matmul` makes this concrete.

```
def tiled_matmul(a: Tensor, b: Tensor, tile_size: int = 64) -> Tensor:
    """Cache-aware matrix multiplication using tiling."""
    # Validate shapes
    if a.shape[-1] != b.shape[-2]:
        raise ValueError(f"Shape mismatch: {a.shape} @ {b.shape}")

    # BLAS libraries automatically implement cache-aware tiling
    # tile_size would control block size in explicit implementation
    result_data = np.matmul(a.data, b.data)
    return Tensor(result_data)
```

: Listing 17.2 — Cache-aware tiled matmul API. `{#lst-17-acceleration-tilled-matmul}`

### 27.6.4 Kernel Fusion

Element-wise operations like GELU are memory-bound: they spend more time moving data than computing on it. Consider the GELU formula:

$$\text{GELU}(x) = 0.5 * x * (1 + \tanh(\sqrt{2/\pi} * (x + 0.044715 * x^3)))$$

A naive implementation creates seven intermediate arrays:

The code in `?@lst-17-acceleration-unfused-gelu` makes this concrete.

```
def unfused_gelu(x: Tensor) -> Tensor:
    """Unfused GELU - creates many temporary arrays."""
    sqrt_2_over_pi = np.sqrt(2.0 / np.pi)

    temp1 = Tensor(x.data**3)                # x³
    temp2 = Tensor(0.044715 * temp1.data)    # 0.044715 * x³
    temp3 = Tensor(x.data + temp2.data)      # x + 0.044715 * x³
    temp4 = Tensor(sqrt_2_over_pi * temp3.data) # √(2/π) * (...)
    temp5 = Tensor(np.tanh(temp4.data))      # tanh(...)
    temp6 = Tensor(1.0 + temp5.data)         # 1 + tanh(...)
    temp7 = Tensor(x.data * temp6.data)      # x * (1 + tanh(...))
    result = Tensor(0.5 * temp7.data)        # 0.5 * x * (...)

    return result
```

: Listing 17.3 — Unfused GELU creating seven intermediate tensors. `{#lst-17-acceleration-unfused-gelu}`

Every temporary writes to memory and the next operation reads it back. For a 4,000,000-element tensor, the unfused version issues 56 million memory operations (7 reads + 7 writes per element). At 50 GB/s — typical desktop CPU bandwidth — moving 214 MB takes 4.48 ms. That’s just the memory traffic, before any actual arithmetic.

Kernel fusion combines all operations into a single expression:

The code in `?@lst-17-acceleration-fused-gelu` makes this concrete.

```
def fused_gelu(x: Tensor) -> Tensor:
    """Fused GELU - all operations in single kernel."""
    sqrt_2_over_pi = np.sqrt(2.0 / np.pi)

    # Single expression - no intermediate arrays
    result_data = 0.5 * x.data * (
        1.0 + np.tanh(sqrt_2_over_pi * (x.data + 0.044715 * x.data**3))
    )

    return Tensor(result_data)
```

: Listing 17.4 — Fused GELU collapsing all ops into a single kernel. `{#lst-17-acceleration-fused-gelu}`

Now there are only two memory operations: read the input, write the output. For the same `{python} fusion_tensor_elements` element tensor, that’s just `{python} fusion_fused_traffic_mb` MB of memory traffic, completing in `{python} fusion_fused_time_ms` milliseconds. The fused version is `{python} fu-`

`sion_speedupx` faster purely from memory bandwidth reduction, even though both versions perform the exact same arithmetic.

**Complexity framing:** both versions perform  $\Theta(N)$  FLOPs on an  $N$ -element tensor — the asymptotic compute class is identical. What fusion changes is the *memory complexity*: the unfused chain is  $\Theta(kN)$  bytes transferred for a  $k$ -op pipeline (one read + one write per intermediate), while the fused kernel is  $\Theta(N)$  bytes (one read of  $x$ , one write of the result). Seven intermediates yields roughly  $7\times$  the traffic. Since element-wise ops are memory-bound, memory complexity — not FLOP count — is what determines wall-clock time. But bandwidth is only half the story when deploying on parallel hardware.

#### **i Systems Implication: Kernel Fusion and Warp Divergence**

On modern GPUs, kernel fusion does more than just save memory bandwidth—it preserves computational harmony. GPUs execute threads in lockstep groups called “warps” (typically 32 threads in NVIDIA architectures). If an unfused series of operations introduces complex branching logic (e.g., the `if x > 0` condition in a ReLU activation), threads within the same warp might evaluate the condition differently, taking divergent execution paths. This causes “warp divergence.” When a warp diverges, the GPU cannot execute the branches simultaneously; it must serialize the execution of each branch while masking out inactive threads, severely degrading hardware utilization. Expertly authored fused kernels are designed to minimize divergent branches and keep data continuously in registers, ensuring the entire warp executes the math in perfect, high-throughput unison without stalling for memory or serialized branches.

### 27.6.5 Parallel Processing

Modern CPUs have multiple cores that can execute operations simultaneously. BLAS libraries automatically spawn threads to parallelize matrix multiplication across cores. A 4-core system can theoretically achieve 4x speedup on compute-bound operations.

However, parallel processing has overhead. Creating threads, synchronizing results, and merging data takes time. For small matrices, this overhead exceeds the benefit. BLAS libraries use heuristics to decide when to parallelize: large matrices get multiple threads, small matrices run on a single core.

This is why real speedups are sublinear. A 4-core system typically lands at 3x rather than 4x:

- Thread creation and destruction overhead
- Cache-coherence traffic between cores
- Memory-bandwidth saturation (all cores share one memory bus)
- Load imbalance (some threads finish before others)

### 27.6.6 Hardware Acceleration

This module uses NumPy and BLAS for CPU acceleration. Production frameworks go further with specialized hardware:

**GPUs** have thousands of simple cores optimized for data parallelism. A matrix multiplication that takes 100ms on a CPU can complete in 1ms on a GPU - a 100x speedup. But GPUs require explicit data transfer between CPU and GPU memory, and this transfer can dominate small operations.

**TPUs** (Tensor Processing Units) are Google’s custom accelerators with systolic array architectures designed specifically for matrix multiplication. A TPU can sustain 100+ TFLOPS on matrix operations.

The acceleration techniques you implement in this module - vectorization, fusion, and cache awareness - apply to all these platforms. The specific implementations differ, but the principles remain constant.

### 27.6.7 Arithmetic Intensity and the Roofline Model

Not all operations are created equal. The roofline model predicts whether an operation will be limited by memory bandwidth or by raw compute. Its driving number is *arithmetic intensity* — floating-point operations per byte transferred:

$$\text{Arithmetic Intensity (AI)} = \text{FLOPs} / \text{Bytes}$$

Element-wise addition of two N-element float32 arrays:

- FLOPs: N (one add per element)
- Bytes:  $3N \times 4 = 12N$  (read A, read B, write C)
- AI =  $1/12 \approx 0.083$  FLOPs/byte

Matrix multiplication of two N×N float32 matrices:

- FLOPs:  $2N^3$  ( $N^3$  multiplies +  $N^3$  adds)
- Bytes:  $3N^2 \times 4 = 12N^2$  (read A, read B, write C)
- AI =  $N/6$  FLOPs/byte

For N=1024 that's 171 FLOPs/byte — about 2048x more arithmetic per byte than element-wise addition. That gap is the whole reason GPUs and tensor cores are spectacular at matmul and uninspiring at element-wise ops.

Table 27.4 classifies each operation by arithmetic intensity and the optimisation strategy that applies.

Table 27.4: Arithmetic intensity and optimization strategy by operation.

| Operation                   | Arithmetic Intensity | Bottleneck         | Optimization Strategy |
|-----------------------------|----------------------|--------------------|-----------------------|
| Element-wise add            | ~0.08 FLOPs/byte     | Memory bandwidth   | Kernel fusion         |
| Element-wise multiply       | ~0.08 FLOPs/byte     | Memory bandwidth   | Kernel fusion         |
| GELU activation             | ~1.0 FLOPs/byte      | Memory bandwidth   | Kernel fusion         |
| Matrix multiply (1024×1024) | ~171 FLOPs/byte      | Compute throughput | Vectorization, tiling |

The roofline model plots achievable performance against arithmetic intensity. Your hardware has a peak memory bandwidth (horizontal line) and peak computational throughput (diagonal line). The minimum of these two lines is your performance ceiling.

## 27.7 Common Errors

These are the errors you'll encounter when optimizing neural networks. Understanding them will save you from subtle performance bugs.

### 27.7.1 Shape Mismatches in Vectorized Code

**Error:** `ValueError: shapes (128, 256) and (128, 512) not aligned`

Matrix multiplication requires inner dimensions to match. For `A @ B`, `A.shape[-1]` must equal `B.shape[-2]`. This error occurs when you try to multiply incompatible shapes.

**Fix:** Always validate shapes before matrix operations:

```
assert a.shape[-1] == b.shape[-2], f"Shape mismatch: {a.shape} @ {b.shape}"
```

### 27.7.2 Memory Bandwidth Bottlenecks

**Symptom:** GPU shows 20% utilization but code is still slow

This indicates a memory-bound operation. The GPU cores are idle, waiting for data from memory. Element-wise operations often hit this bottleneck.

**Fix:** Use kernel fusion to reduce memory traffic. Combine multiple element-wise operations into a single fused kernel.

### 27.7.3 Cache Thrashing

**Symptom:** Performance degrades dramatically for matrices larger than  $1024 \times 1024$

When your working set exceeds cache size, the CPU spends most of its time loading data from main memory rather than computing.

**Fix:** Use tiling/blocking to keep working sets in cache. Break large matrices into smaller tiles that fit in L2 or L3 cache.

### 27.7.4 False Dependencies

**Symptom:** Parallel code runs slower than sequential code

Creating temporary arrays in a loop can prevent parallelization because each iteration depends on the previous one's memory allocation.

**Fix:** Preallocate output arrays and reuse them:

```
# Bad: creates new array each iteration
for i in range(1000):
    result = x + y

# Good: reuses same output array
result = np.zeros_like(x)
for i in range(1000):
    np.add(x, y, out=result)
```

## 27.8 Production Context

### 27.8.1 Your Implementation vs. PyTorch

Your acceleration techniques demonstrate the same principles PyTorch uses internally. The difference is scale: PyTorch supports GPUs, automatic kernel fusion through compilers, and thousands of optimized operations.

Table 27.5 places your implementation side by side with the production reference for direct comparison.

Table 27.5: Feature comparison between TinyTorch acceleration and PyTorch internals.

| Feature              | Your Implementation | PyTorch             |
|----------------------|---------------------|---------------------|
| <b>Vectorization</b> | NumPy BLAS          | CUDA/cuBLAS for GPU |

| Feature                | Your Implementation  | PyTorch                       |
|------------------------|----------------------|-------------------------------|
| <b>Kernel Fusion</b>   | Manual fusion        | Automatic via TorchScript/JIT |
| <b>Backend</b>         | CPU only             | CPU, CUDA, Metal, ROCm        |
| <b>Multi-threading</b> | Automatic (OpenBLAS) | Configurable thread pools     |
| <b>Operations</b>      | ~5 accelerated ops   | 2000+ optimized ops           |

### 27.8.2 Code Comparison

The following comparison shows how acceleration appears in TinyTorch versus PyTorch. The API patterns are similar, but PyTorch adds GPU support and automatic optimization.

## 27.9 Your TinyTorch

```
from tinytorch.perf.acceleration import vectorized_matmul, fused_gelu

# CPU-based acceleration
x = Tensor(np.random.randn(128, 512))
w = Tensor(np.random.randn(512, 256))

# Vectorized matrix multiplication
h = vectorized_matmul(x, w)

# Fused activation
output = fused_gelu(h)
```

## 27.10 PyTorch

```
import torch

# GPU acceleration with same concepts
x = torch.randn(128, 512, device='cuda')
w = torch.randn(512, 256, device='cuda')

# Vectorized (cuBLAS on GPU)
h = x @ w

# Fused via JIT compilation
@torch.jit.script
def fused_gelu(x):
    return 0.5 * x * (1 + torch.tanh(0.797885 * (x + 0.044715 * x**3)))

output = fused_gelu(h)
```

Let's walk through the key differences:

- **Line 1 (Import):** TinyTorch provides explicit acceleration functions; PyTorch integrates acceleration into the core tensor operations.
- **Line 4-5 (Device):** TinyTorch runs on CPU via NumPy; PyTorch supports `device='cuda'` for GPU acceleration.
- **Line 8 (Matrix multiply):** Both use optimized BLAS, but PyTorch uses cuBLAS on GPU for 10-100x additional speedup.
- **Line 11-13 (Fusion):** TinyTorch requires manual fusion; PyTorch's JIT compiler can automatically fuse operations.
- **Performance:** For this example, TinyTorch might take 5ms on CPU; PyTorch takes 0.05ms on GPU - a 100x speedup.

#### 💡 What's Identical

The acceleration principles: vectorization reduces instruction count, fusion reduces memory traffic, and hardware awareness guides optimization choices. These concepts apply everywhere.

### 27.10.1 Why Acceleration Matters at Scale

Three numbers explain why every framework team obsesses over kernel performance:

- **GPT-3 training:** 175B parameters  $\times$  300B tokens  $\approx 10^{23}$  FLOPs. Naive code: centuries. Optimized TPUs: weeks.
- **Real-time inference:** 1000 requests/second forces **sub-millisecond latency** per request. Every 2x speedup doubles throughput at the same dollar cost.
- **Cost efficiency:** cloud GPU time runs \$2–10/hour. A 2x speedup saves **\$1000–5000/week** on a single production model.

At this scale, the percent improvements you ship in this module compound into millions in savings — and into capabilities that simply weren't reachable on the slow path.

## 27.11 Check Your Understanding

### 💡 Check Your Understanding — Acceleration

Before moving on, verify you can articulate each of the following:

- What kernel fusion actually fuses, and how warp divergence on GPUs can erase the memory-bandwidth gain.
- Why matrix multiplication sits at ~171 FLOPs/byte (compute-bound) while GELU sits at ~1 FLOPs/byte (memory-bound), and what each regime demands.
- How BLAS Level 1/2/3 operations differ in arithmetic intensity, and why every ML framework ultimately routes through GEMM.
- Why cache-aware tiling matters once the working set exceeds L2, and how to pick a tile size from the cache budget.

If any of these feels fuzzy, revisit the Core Concepts section (especially Kernel Fusion and Arithmetic Intensity) before moving on.

Test your understanding of acceleration techniques with these quantitative questions.

#### Q1: Arithmetic Intensity

Matrix multiplication of two  $1024 \times 1024$  float32 matrices performs 2,147,483,648 FLOPs. It reads 4 MB (A) + 4 MB (B) and writes 4 MB (C), so 12 MB of memory traffic total. What is the arithmetic intensity?

#### 💡 Answer

Arithmetic Intensity = 2,147,483,648 FLOPs / 12,582,912 bytes = **~171 FLOPs/byte**

This high arithmetic intensity — compared to ~0.08 for element-wise ops — is why matrix multiplication is ideal for GPUs and why it dominates neural network training time.

#### Q2: Memory Bandwidth Savings

Your fused GELU processes a tensor with 1,000,000 elements (~4 MB as float32). The unfused version creates 7 intermediate arrays. How much memory bandwidth does fusion save?

#### 💡 Answer

**Unfused:** 7 reads + 7 writes + 1 input read + 1 output write = 16 ops  $\times$  ~4 MB  $\approx$  **61 MB**

**Fused:** 1 input read + 1 output write = 2 ops  $\times$  ~4 MB  $\approx$  **8 MB**

**Savings:** 61 – 8  $\approx$  **53 MB saved (~87.5% reduction)**

At ~50 GB/s CPU bandwidth that's roughly 1 ms saved per GELU call. A transformer with 96 GELU activations per forward pass recovers ~96 ms — a 10–20% throughput win on inference for free.

#### Q3: Cache Tiling

A CPU has 256 KB L2 cache. You're multiplying two  $2048 \times 2048$  float32 matrices (16 MB each). What tile size keeps the working set in L2 cache?

#### 💡 Answer

Tiled multiplication needs three tiles resident at once:

- Tile from A:  $\text{tile\_size} \times \text{tile\_size} \times 4$  bytes
- Tile from B:  $\text{tile\_size} \times \text{tile\_size} \times 4$  bytes
- Output tile:  $\text{tile\_size} \times \text{tile\_size} \times 4$  bytes

Constraint:  $3 \times \text{tile\_size}^2 \times 4 \leq 256 \text{ KB}$

Solving:  $\text{tile\_size}^2 \leq 262,144 / 12 \approx 21,845$ , so **tile\_size  $\approx$  147**.

In practice, snap to a power of two: **128 works well** ( $3 \times 128^2 \times 4 = 192 \text{ KB}$ , leaving headroom for other data).

#### Q4: BLAS Performance

Your vectorized matmul completes a  $1024 \times 1024$  multiplication in 10 ms. The operation requires 2.15 billion FLOPs. What is your achieved performance in GFLOPS?

#### 💡 Answer

GFLOPS = 2,150,000,000 FLOPs / (0.01 s  $\times$  1,000,000,000) = **215 GFLOPS**

For reference:

- Modern CPU peak: 500–1000 GFLOPS (AVX-512)
- Your efficiency: 215 / 500 = **~43% of peak** (typical for real code)
- GPU equivalent: ~50 TFLOPS (about 230x faster than a single CPU core)

#### Q5: Speedup from Fusion

Unfused GELU takes 8 ms on a  $2000 \times 2000$  tensor. Fused GELU takes 2.5 ms. What percentage of the unfused time was memory overhead?



 Answer

Speedup = 8 ms / 2.5 ms = **3.2x faster**

Both versions perform the same arithmetic, so the gap is memory bandwidth:

- Memory overhead =  $(8 - 2.5) / 8 = 68.75\%$

Nearly **69% of the unfused version's runtime was spent waiting for memory**. That's typical for element-wise operations with low arithmetic intensity.

## 27.12 Key Takeaways

- **Arithmetic intensity decides the ceiling:** FLOPs-per-byte places every operation on the roofline and tells you whether vectorization (compute-bound) or fusion (memory-bound) is the right lever.
- **Fusion changes memory complexity, not FLOP complexity:**  $\Theta(kN)$  bytes collapses to  $\Theta(N)$  bytes when  $k$  element-wise ops are fused — identical arithmetic, dramatically less traffic.
- **GEMM dominates because it is compute-bound:**  $\sim 171$  FLOPs/byte at  $N=1024$  is why decades of hardware design revolve around matrix multiplication and why BLAS Level-3 is the hot path in every framework.
- **Hardware-aware, not clever:** acceleration rarely needs a new algorithm — it needs code that stops fighting the cache, the SIMD lanes, and the warp scheduler.

**Coming next:** Module 18 closes the loop by eliminating the call entirely — KV caching turns autoregressive generation from  $O(n^2)$  to  $O(n)$  work by memoizing past keys and values.

## 27.13 Further Reading

For students who want to understand the academic foundations and implementation details of neural network acceleration:

### 27.13.1 Seminal Papers

- **Roofline Model** - Williams et al. (2009). The foundational framework for understanding performance bottlenecks based on arithmetic intensity. Essential for diagnosing whether your code is compute-bound or memory-bound. [IEEE](#)
- **BLAS: The Basic Linear Algebra Subprograms** - Lawson et al. (1979). The specification that defines standard matrix operations. Every ML framework ultimately calls BLAS for performance-critical operations. **Systems Implication:** Defined the memory hierarchy abstractions (registers, L1/L2 cache, RAM) allowing matrix multiplication to be tiled for optimal cache reuse, fundamentally defining modern dense compute limits. [ACM TOMS](#)
- **Optimizing Matrix Multiplication** - Goto & Geijn (2008). Detailed explanation of cache blocking, register tiling, and microkernel design for high-performance GEMM. This is how BLAS libraries achieve near-peak performance. [ACM TOMS](#)
- **TVM: An Automated End-to-End Optimizing Compiler** - Chen et al. (2018). Demonstrates automatic optimization including kernel fusion and memory planning for deep learning. Shows how compilers can automatically apply the techniques you learned manually. [OSDI](#)

### 27.13.2 Additional Resources

- **Tutorial:** [“What Every Programmer Should Know About Memory”](#) by Ulrich Drepper - Deep dive into cache hierarchies and their performance implications

- **Documentation:** [Intel MKL Developer Reference](#) - See how production BLAS libraries implement vectorization and threading

## 27.14 What's Next

You just made each operation as cheap as the hardware allows. The next move is to skip operations entirely.

### Coming Up: Module 18 - Memoization

Acceleration shrinks the cost of every call. **Memoization eliminates the call.** Module 18 builds KV-caching for transformer generation, caches repeated forward passes, and reuses attention patterns so the model never recomputes what it already knows. The two techniques compose: a fused kernel that gets called zero times is the cheapest kernel of all.

### Preview - How Acceleration Gets Used in Future Modules:

Table 27.6 traces how this module is reused by later parts of the curriculum.

Table 27.6: How acceleration composes with memoization, benchmarking, and capstone.

| Module                  | What It Does                       | Your Acceleration In Action                                      |
|-------------------------|------------------------------------|------------------------------------------------------------------|
| <b>18: Memoization</b>  | Cache repeated computations        | Fused kernels + KV cache minimize memory traffic                 |
| <b>19: Benchmarking</b> | Systematic performance measurement | <code>benchmark(vectorized_matmul, sizes=[128, 256, 512])</code> |
| <b>20: Capstone</b>     | Complete optimized model           | Acceleration throughout model pipeline                           |

## 27.15 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.

### Performance Note

Acceleration techniques depend on hardware. Results will vary between CPUs. Use Module 14's profiler to measure your specific hardware's characteristics.

## 🔥 Chapter 28

# Module 18: Memoization

Autoregressive inference recomputes the same attention keys and values for every new token, billions of times across a production fleet. The KV cache is the memoization trick that turns that  $O(N^2)$  redundancy into  $O(N)$  by trading GPU HBM for saved compute. This module builds the cache, and the fragmentation problem at long context that makes production systems reach for PagedAttention.

### **i** Module Info

**OPTIMIZATION TIER** | Difficulty: ●●○○ | Time: 3-5 hours | Prerequisites: 01-14

**Prerequisites:** **Modules 01-14** means you should be comfortable with:

- Tensor operations, matrix multiplication, and shape manipulation (Module 01)
- Transformer architectures and attention (Modules 12-13)
- Profiling tools (Module 14) to measure speedup

If you understand how transformers compute attention and why it's expensive, you're ready to learn how to make inference dramatically faster.

## 28.1 Overview

When ChatGPT writes a 100-word response, the naive transformer would recompute the keys and values for every previous token at every step, doing 5,050 K,V projections to produce 100 tokens. Almost all of that work is duplicate. Memoize it once and the cost collapses to 100.

That single change is what makes real-time chat affordable. KV caching stores the key and value matrices from past tokens so each new token only computes its own — turning generation from  $O(n^2)$  into  $O(n)$  and unlocking the 10–15x speedup every deployed LLM relies on.

In this module you build that cache. By the end you will have a `KVCache` class with  $O(1)$  updates, a non-invasive hook that retrofits caching onto an existing transformer, and a quantitative grasp of the memory-versus-compute trade you just bought.

## 28.2 Where Memoization Fits

Acceleration (Module 17) made *any* computation faster — vectorization, cache locality, kernel fusion. Those tricks apply to matmul, convolution, attention, everything.

Memoization is the opposite move: a single, narrow optimization that targets one structural property of autoregressive generation — that past keys and values never change. You spend  $O(n)$  memory to skip  $O(n^2)$  work. It only helps decoder transformers, but on those it is the difference between economically viable and not.

## 28.3 Learning Objectives

💡 By completing this module, you will:

- **Implement** a KVCache class with efficient memory management and  $O(1)$  update operations
- **Master** the memory-compute trade-off: accepting  $O(n)$  memory overhead for  $O(n^2)$  to  $O(n)$  speedup
- **Understand** why memoization transforms generation complexity from quadratic to linear
- **Connect** your implementation to production systems like ChatGPT and Claude that rely on KV caching

## 28.4 What You'll Build




Figure 28.1: **TinyTorch KV Caching**: Reusing previous computations to speed up text generation.

### Implementation roadmap:

Table 28.1 lays out the implementation in order, one part at a time.

Table 28.1: **Implementation roadmap for the KV cache data structure.**

| Step | What You'll Implement           | Key Concept                           |
|------|---------------------------------|---------------------------------------|
| 1    | <code>KVCache.__init__()</code> | Pre-allocated cache storage per layer |
| 2    | <code>KVCache.update()</code>   | $O(1)$ cache append without copying   |
| 3    | <code>KVCache.get()</code>      | $O(1)$ retrieval of cached values     |
| 4    | <code>enable_kv_cache()</code>  | Non-invasive model enhancement        |
| 5    | Performance analysis            | Measure speedup and memory usage      |

### The pattern you'll enable:

```

# Enable caching for dramatic speedup
cache = enable_kv_cache(model)
# Generate with 10-15x faster inference
  
```

```
output = model.generate(prompt, max_length=100)
```

### 28.4.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Multi-batch cache management (production systems handle thousands of concurrent sequences)
- Cache eviction strategies (handling sequences longer than `max_seq_len`)
- GPU memory optimization (production uses memory pools and paging)
- Speculative decoding (advanced technique that builds on KV caching)

**You are building the core memoization mechanism.** Advanced cache management comes in production deployment.

## 28.5 API Reference

This section provides a quick reference for the `KVCache` class you'll build. Use this as your guide while implementing and debugging.

### 28.5.1 KVCache Constructor

```
KVCache(batch_size: int, max_seq_len: int, num_layers: int,
        num_heads: int, head_dim: int) -> KVCache
```

Pre-allocates cache storage for all transformer layers. Each layer gets two tensors (keys and values) sized to hold the maximum sequence length.

**Parameters:** - `batch_size`: Number of sequences to cache simultaneously - `max_seq_len`: Maximum sequence length to support - `num_layers`: Number of transformer layers in the model - `num_heads`: Number of attention heads per layer - `head_dim`: Dimension of each attention head

### 28.5.2 Core Methods

Table 28.2 lists the methods and helpers you will implement.

Table 28.2: Core methods on the `KVCache` class.

| Method                        | Signature                                                                  | Description                                           |
|-------------------------------|----------------------------------------------------------------------------|-------------------------------------------------------|
| <code>update</code>           | <code>update(layer_idx: int, key: Tensor, value: Tensor) -&gt; None</code> | Append new K,V to cache for given layer               |
| <code>get</code>              | <code>get(layer_idx: int) -&gt; Tuple[Tensor, Tensor]</code>               | Retrieve cached K,V for attention computation         |
| <code>advance</code>          | <code>advance() -&gt; None</code>                                          | Move sequence position forward after processing token |
| <code>reset</code>            | <code>reset() -&gt; None</code>                                            | Clear cache for new generation sequence               |
| <code>get_memory_usage</code> | <code>get_memory_usage() -&gt; Dict[str, float]</code>                     | Calculate cache memory consumption                    |

### 28.5.3 Helper Functions

Table 28.3 lists the methods and helpers you will implement.

Table 28.3: Helper functions for enabling and disabling the KV cache.

| Function                      | Signature                                         | Description                               |
|-------------------------------|---------------------------------------------------|-------------------------------------------|
| <code>enable_kv_cache</code>  | <code>enable_kv_cache(model) -&gt; KVCache</code> | Non-invasively add caching to transformer |
| <code>disable_kv_cache</code> | <code>disable_kv_cache(model) -&gt; None</code>   | Restore original attention behavior       |

## 28.6 Core Concepts

This section covers the fundamental ideas you need to understand memoization in transformers. These concepts explain why KV caching is the optimization that makes production language models economically viable.

### 28.6.1 Caching Computation

Memoization trades memory for speed by storing computation results for reuse. When a function is called with the same inputs repeatedly, computing the result once and caching it eliminates redundant work. This trade-off makes sense when memory is cheaper than computation, which is almost always true for inference.

In transformers, attention is the perfect target for memoization. During autoregressive generation, each new token requires attention over all previous tokens. The naive approach recomputes key and value projections for every previous token at every step, leading to quadratic complexity. But these projections never change once computed, making them ideal candidates for caching.

Here's the core insight in your implementation:

The code in `?@lst-18-memoization-update` makes this concrete.

```
def update(self, layer_idx: int, key: Tensor, value: Tensor) -> None:
    """Update cache with new key-value pairs for given layer."""
    if layer_idx >= self.num_layers:
        raise ValueError(f"Layer index {layer_idx} >= num_layers {self.num_layers}")

    if self.seq_pos >= self.max_seq_len:
        raise ValueError(f"Sequence position {self.seq_pos} >= max_seq_len {self.max_seq_len}")

    # Get cache for this layer
    key_cache, value_cache = self.caches[layer_idx]

    # Update cache at current position (efficient O(1) write)
    key_cache.data[:, :, self.seq_pos:self.seq_pos+1, :] = key.data
    value_cache.data[:, :, self.seq_pos:self.seq_pos+1, :] = value.data
```

: Listing 18.1 — O(1) KV cache append into pre-allocated per-layer storage. `{#lst-18-memoization-update}`

This  $O(1)$  update operation writes directly to a pre-allocated position in the cache. No array resizing, no data copying, just an indexed assignment. The use of `.data` accesses the underlying NumPy array directly, avoiding gradient tracking overhead since caching is inference-only.

The computational savings compound across generation steps. For a 100-token sequence:

- Without caching:  $1 + 2 + 3 + \dots + 100 = 5,050$  K,V computations
- With caching: 100 K,V computations (one per token)
- Speedup: 50x reduction in K,V computation alone

### 28.6.2 KV Cache in Transformers

Transformer attention computes three projections from the input: query (Q), key (K), and value (V). The attention output is computed as  $\text{softmax}(Q @ K^T / \sqrt{d_k}) @ V$ . During generation, each new token produces a new query, but the keys and values from previous tokens remain constant.

Consider generating the sequence “Hello world!”:

```
Step 1: Input = ["Hello"]
        Compute: Q, K, V
        Attention: Q @ [K] @ [V]

Step 2: Input = ["Hello", "world"]
        Compute: Q, K, V
        Attention: Q @ [K, K] @ [V, V]
        Problem: K and V are recomputed unnecessarily!

Step 3: Input = ["Hello", "world", "!"]
        Compute: Q, K, V
        Attention: Q @ [K, K, K] @ [V, V, V]
        Problem: K, V, K, V are all recomputed!
```

The cache eliminates this redundancy:

```
Step 1: Compute K, V → Cache them
Step 2: Compute K, V → Append to cache
        Attention: Q @ cached[K, K] @ cached[V, V]
Step 3: Compute K, V → Append to cache
        Attention: Q @ cached[K, K, K] @ cached[V, V, V]
```

Each step now computes only one new K,V pair instead of recomputing all previous pairs. This algorithmic optimization dramatically accelerates single-batch inference, but scaling this to thousands of concurrent users introduces a formidable systems engineering challenge.

#### **i** Systems Implication: KV-Cache Memory Fragmentation and PagedAttention

While KV caching transforms computational time complexity from  $O(n^2)$  to  $O(n)$ , it creates a severe systems bottleneck: memory fragmentation. Because the final length of an autoregressively generated sequence is unknown ahead of time, naive inference engines must statically pre-allocate maximum-length, contiguous memory blocks for every request's cache. If a request finishes early, the remaining allocated memory is wasted, leading to massive internal fragmentation. Modern production engines like vLLM solve this with **PagedAttention**. Inspired by operating system virtual memory, PagedAtten-

tion divides the KV-cache into fixed-size, non-contiguous “pages.” By dynamically allocating pages on demand and mapping them to a virtual sequence space, this approach nearly eliminates internal fragmentation, allowing a single GPU to serve exponentially larger batch sizes and drastically reducing the cost per token.

### 28.6.3 Gradient Checkpointing

While KV caching optimizes inference, gradient checkpointing addresses the opposite problem: memory consumption during training. Training requires storing intermediate activations for backpropagation, but for very deep networks, this can exceed available memory. Gradient checkpointing trades compute for memory by not storing all activations.

The technique works by discarding some intermediate activations during the forward pass and recomputing them during backpropagation when needed. Instead of storing activations for all layers (requiring  $O(n)$  memory where  $n$  is the number of layers), checkpointing only stores activations at regular intervals (checkpoints). Between checkpoints, activations are recomputed from the last checkpoint during the backward pass.

For a transformer with 96 layers:

- Without checkpointing: Store 96 sets of activations
- With checkpointing every 12 layers: Store 8 sets, recompute 11 sets during backward
- Memory reduction: 12x decrease
- Compute increase: ~33% slower training (recomputation overhead)

This is the inverse trade-off from KV caching. KV caching spends memory to save compute during inference. Gradient checkpointing spends compute to save memory during training. Both techniques recognize that memory and compute are fungible resources with different costs in different contexts.

### 28.6.4 Cache Invalidation

Cache invalidation is famously hard. Autoregressive generation makes it easy: each cached K,V pair is valid for the entire sequence being generated, and the entire cache is thrown away the moment a new sequence starts.

That simplicity falls out of the autoregressive property — every token depends only on tokens that came before it, and once those dependencies are computed they never change.

Here’s how your implementation handles cache lifecycle:

The code in `?@lst-18-memoization-reset` makes this concrete.

```
def reset(self) -> None:
    """Reset cache for new generation sequence."""
    self.seq_pos = 0

    # Zero out caches for clean state (helps with debugging)
    for layer_idx in range(self.num_layers):
        key_cache, value_cache = self.caches[layer_idx]
        key_cache.data.fill(0.0)
        value_cache.data.fill(0.0)
```

: Listing 18.2 — Cache reset clearing per-layer K and V storage between sequences. {#lst-18-memoization-reset}



The reset operation returns the sequence position to zero and clears the cache data. This is called when starting to generate a new sequence, ensuring no stale data from previous generations affects the current one.

Production systems handle more complex invalidation scenarios:

- **Max length reached:** When the sequence fills the cache, either error out or implement a sliding window
- **Batch inference:** Each sequence in a batch has independent cache state
- **Multi-turn conversation:** Some systems maintain cache across turns, others reset per turn

### 28.6.5 Memory-Compute Trade-offs

Every optimization involves trade-offs. KV caching trades memory for speed, and understanding this exchange quantitatively reveals when the technique makes sense.

For a transformer with  $L$  layers,  $H$  heads per layer, dimension  $D$  per head, and maximum sequence length  $S$ , the cache requires:

$$\text{Memory} = 2 \times L \times H \times S \times D \times 4 \text{ bytes}$$

Example (GPT-2 Small):

$L = 12$  layers,  $H = 12$  heads,  $S = 1024$  tokens,  $D = 64$  dims

$$\text{Memory} = 2 \times 12 \times 12 \times 1024 \times 64 \times 4 = 75,497,472 \text{ bytes} \approx 72 \text{ MB}$$

For a 125M-parameter model (500 MB of weights), that cache adds 14% memory overhead — modest, until you weigh it against the compute it saves.

Without caching, generating a sequence of length  $N$  requires computing  $K, V$  for:

- Step 1: 1 token
- Step 2: 2 tokens
- Step 3: 3 tokens
- Step  $N$ :  $N$  tokens
- Total:  $1 + 2 + 3 + \dots + N = N(N+1)/2 \approx N^2/2$  computations

With caching:

- Step 1: 1 token (compute and cache)
- Step 2: 1 token (compute and append)
- Step 3: 1 token (compute and append)
- Step  $N$ : 1 token (compute and append)
- Total:  $N$  computations

For  $N = 100$  tokens, caching gives a 50x reduction in  $K, V$  computation. At  $N = 1000$ , the reduction is 500x. The savings grow with sequence length, which is exactly why long-context models depend on this trick.

Table 28.4 shows how KV-cache savings scale with sequence length.

Table 28.4: KV-cache memory and compute savings as sequence length grows.

| Sequence Length | Cache Memory | Compute Reduction | Effective Speedup |
|-----------------|--------------|-------------------|-------------------|
| 10 tokens       | 72 MB        | 5.5x              | 3-5x              |
| 50 tokens       | 72 MB        | 25.5x             | 8-12x             |
| 100 tokens      | 72 MB        | 50.5x             | 10-15x            |
| 500 tokens      | 72 MB        | 250.5x            | 12-20x            |

The effective speedup is lower than the theoretical compute reduction because attention includes other operations beyond K,V projection, but the benefit is still dramatic.

## 28.7 Common Errors

These are the errors you'll encounter most often when implementing KV caching. Understanding why they happen will save hours of debugging.

### 28.7.1 Cache Position Out of Bounds

**Error:** `ValueError: Sequence position 128 >= max_seq_len 128`

This happens when you try to append to a full cache. The cache is pre-allocated with a maximum sequence length, and attempting to write beyond that length raises an error.

**Cause:** Generation exceeded the maximum sequence length specified when creating the cache.

**Fix:** Either increase `max_seq_len` when creating the cache, or implement cache eviction logic to handle sequences longer than the maximum.

```
# Create cache with sufficient capacity
cache = KVCache(batch_size=1, max_seq_len=2048, # Increased from 128
                num_layers=12, num_heads=12, head_dim=64)
```

### 28.7.2 Forgetting to Advance Position

**Error:** Cache retrieval returns the same K,V repeatedly, or update overwrites previous values

**Symptom:** Generated text repeats, or cache doesn't grow as expected

**Cause:** Forgetting to call `cache.advance()` after updating all layers for a token.

**Fix:** Always advance the cache position after processing a complete token through all layers:

```
for layer_idx in range(num_layers):
    cache.update(layer_idx, new_key, new_value)

cache.advance() # Move to next position for next token
```

### 28.7.3 Shape Mismatches

**Error:** Broadcasting error or shape mismatch when updating cache

**Symptom:** `ValueError: could not broadcast input array from shape (1,8,64,64) into shape (1,8,1,64)`

**Cause:** The key and value tensors passed to `update()` must have shape `(batch, heads, 1, head_dim)` with sequence dimension equal to 1 (single new token).

**Fix:** Ensure new K,V tensors represent a single token:

```
# Correct: Single token (seq_len = 1)
new_key = Tensor(np.random.randn(batch_size, num_heads, 1, head_dim))
cache.update(layer_idx, new_key, new_value)

# Wrong: Multiple tokens (seq_len = 64)
wrong_key = Tensor(np.random.randn(batch_size, num_heads, 64, head_dim))
cache.update(layer_idx, wrong_key, wrong_value) # This will fail!
```

### 28.7.4 Cache Not Reset Between Sequences

**Error:** Second generation includes tokens from first generation

**Symptom:** Model generates text that seems to continue from a previous unrelated sequence

**Cause:** Forgetting to reset the cache when starting a new generation sequence.

**Fix:** Always reset the cache before generating a new sequence:

```
# Generate first sequence
output1 = model.generate(prompt1)

# Reset cache before second sequence
model._kv_cache.reset()

# Generate second sequence (independent of first)
output2 = model.generate(prompt2)
```

## 28.8 Production Context

### 28.8.1 Your Implementation vs. PyTorch

Your KVCache implementation uses the same conceptual design as production frameworks. The differences lie in scale, optimization level, and integration depth. PyTorch's KV cache implementation is written in C++ and CUDA for speed, supports dynamic batching for serving multiple users, and includes sophisticated memory management with paging and eviction.

Table 28.5 places your implementation side by side with the production reference for direct comparison.

Table 28.5: Feature comparison between TinyTorch KV cache and PyTorch transformers.

| Feature                  | Your Implementation | PyTorch (Transformers library) |
|--------------------------|---------------------|--------------------------------|
| <b>Backend</b>           | NumPy (CPU)         | C++/CUDA (GPU)                 |
| <b>Pre-allocation</b>    | Fixed max_seq_len   | Dynamic growth + paging        |
| <b>Batch support</b>     | Single batch size   | Dynamic batching               |
| <b>Memory management</b> | Simple reset        | LRU eviction, memory pools     |
| <b>Update complexity</b> | O(1)                | O(1) with optimized kernels    |

### 28.8.2 Code Comparison

The following comparison shows how KV caching is used in TinyTorch versus production PyTorch. The API patterns are similar because the underlying concept is identical.

## 28.9 Your TinyTorch

```
from tinytorch.perf.memoization import enable_kv_cache

# Enable caching
cache = enable_kv_cache(model)
```

```
# Generate with caching (10-15x faster)
for _ in range(100):
    logits = model.forward(input_token)
    next_token = sample(logits)
    # Cache automatically used and updated
    input_token = next_token

# Reset for new sequence
cache.reset()
```

## 28.10 PyTorch

```
from transformers import AutoModelForCausalLM

model = AutoModelForCausalLM.from_pretrained("gpt2")

# KV cache enabled automatically during generate()
outputs = model.generate(
    input_ids,
    max_length=100,
    use_cache=True # KV caching enabled
)

# Cache managed internally by HuggingFace
# Automatically reset between generate() calls
```

Let's examine each approach to understand the similarities and differences:

- **Line 1-2 (Imports):** TinyTorch uses an explicit `enable_kv_cache()` function to opt in to caching. PyTorch's Transformers library integrates caching directly into the model architecture.
- **Line 4-5 (Setup):** TinyTorch requires manually enabling the cache and storing the reference. PyTorch handles this transparently when you call `generate()`.
- **Line 7-12 (Generation):** TinyTorch's loop explicitly manages token generation with the cache working behind the scenes. PyTorch's `generate()` method encapsulates the entire loop and automatically uses caching when `use_cache=True`.
- **Line 14-15 (Reset):** TinyTorch requires manual cache reset between sequences. PyTorch automatically resets the cache at the start of each `generate()` call.

The core difference is abstraction level. TinyTorch exposes the cache as an explicit object you control, making the optimization visible for learning. PyTorch hides caching inside `generate()` for ease of use in production. Both implementations use the same  $O(1)$  append pattern you built.

### 💡 What's Identical

The fundamental algorithm: compute K,V once, append to cache, retrieve for attention. Production systems add memory management and batching, but the core optimization is exactly what you implemented.

### 28.10.1 Why Memoization Matters at Scale

To appreciate the production impact of KV caching, consider the economics of language model serving:

- **ChatGPT** serves millions of requests per day. Without KV caching, serving costs would be roughly 10x higher — enough to break the pricing model.
- **GitHub Copilot** generates completions in real time. Without caching, latency would jump from ~100 ms to 1–2 seconds, killing the inline-completion experience.
- **API serving**: a single V100 hosting GPT-2 handles 50–100 concurrent users with caching, but only 5–10 without it. That 10x gap is what determines whether a deployment is profitable.

The memory cost is modest compared to the benefit. For a GPT-2 model:

- Model parameters: 500 MB (loaded once, shared across all users)
- KV cache per user: 75 MB
- 10 concurrent users: 750 MB cache + 500 MB model = 1.25 GB total
- Fits comfortably on a 16 GB GPU while delivering 10x throughput

## 28.11 Check Your Understanding

### 💡 Check Your Understanding — Memoization

Before moving on, verify you can articulate each of the following:

- Why the KV cache dominates inference memory at long context, and how PagedAttention fragments it to reclaim wasted capacity.
- How caching transforms autoregressive generation from  $O(n^2)$  to  $O(n)$  K,V projections, and why the savings compound with sequence length.
- The memory-compute trade-off: when trading  $O(n)$  memory for an  $O(n^2) \rightarrow O(n)$  compute reduction is a good deal (almost always, for inference).
- How cache lifecycle (update, advance, reset) interacts with batched and multi-turn serving.

If any of these feels fuzzy, revisit the Core Concepts section (especially KV Cache in Transformers and Memory-Compute Trade-offs) before moving on.

Test yourself with these systems thinking questions. They're designed to build intuition for the performance characteristics and trade-offs you'll encounter in production ML systems.

#### Q1: Cache Memory Calculation

A 12-layer transformer has 8 attention heads per layer, each head has 64 dimensions, maximum sequence length is 1024, and batch size is 4. Calculate the KV cache memory requirement.

### 💡 Answer

Shape per cache tensor: (batch=4, heads=8, seq=1024, dim=64)

Elements per tensor:  $4 \times 8 \times 1024 \times 64 = 2,097,152$

Each layer has 2 tensors (K and V):  $2 \times 2,097,152 = 4,194,304$  elements per layer

Total across 12 layers:  $12 \times 4,194,304 = 50,331,648$  elements

Memory:  $50,331,648 \times 4 \text{ bytes} = 201,326,592 \text{ bytes} \approx \mathbf{192 \text{ MB}}$

This is why production systems carefully tune batch size and sequence length.

#### Q2: Complexity Reduction

Without caching, generating 200 tokens requires how many K,V computations? With caching?

 **Answer**

**Without caching:**  $1 + 2 + 3 + \dots + 200 = 200 \times 201 / 2 = 20,100$  computations

**With caching:** 200 computations (one per token)

**Reduction:**  $20,100 / 200 = 100.5\times$  fewer K,V computations

This is why the speedup grows with sequence length.

**Q3: Memory-Compute Trade-off**

A model uses 2 GB for parameters. Adding KV cache uses 300 MB. Is this trade-off worthwhile if it provides 12x speedup?

 **Answer**

**Memory overhead:**  $300 \text{ MB} / 2000 \text{ MB} = 15\%$  increase

**Speedup:** 12x faster generation

**Analysis:**

- Cost: 15% more memory
- Benefit: 12x more throughput (or 12x lower latency)
- Result: You can serve 12x more users with 1.15x the memory

**Verdict:** Absolutely worthwhile. Memory is cheap; compute is expensive.

In production, this enables serving 120 users per GPU instead of 10, dramatically reducing infrastructure costs.

**Q4: Cache Hit Rate**

During generation, what percentage of K,V retrievals come from cache vs. fresh computation after 50 tokens?

 **Answer**

At token position 50:

- Fresh computation: 1 new K,V pair
- Cache retrievals: 49 previous K,V pairs
- Total: 50 K,V pairs needed

**Cache hit rate:**  $49/50 = 98\%$

As generation continues:

- Token 100:  $99/100 = 99\%$  hit rate
- Token 500:  $499/500 = 99.8\%$  hit rate

The cache hit rate approaches 100% for long sequences, explaining why speedup increases with length.

**Q5: Batch Inference Scaling**

Cache memory for `batch_size=1` is 75 MB. What is cache memory for `batch_size=8`?

 **Answer**

Cache memory scales linearly with batch size:

**`batch_size=8`:**  $75 \text{ MB} \times 8 = 600 \text{ MB}$

This is why production systems carefully manage batch size:

- Larger batches → higher throughput (more sequences per second)
- Larger batches → more memory (may hit GPU limits)

Trade-off example on 16 GB GPU:

- Model: 2 GB
- Available for cache: 14 GB
- Max batch size: 14 GB / 75 MB  $\approx$  186 sequences

Production systems balance batch size against latency requirements and memory constraints.

## 28.12 Key Takeaways

- **Memoization exploits invariance:** past keys and values never change during autoregressive decoding, so computing them once instead of  $n^2$  times is pure arbitrage.
- **$O(n)$  memory buys an  $O(n^2) \rightarrow O(n)$  compute reduction:** at  $n = 100$  tokens that's a  $50\times$  reduction in K,V work; at  $n = 1000$  it's  $500\times$ .
- **Memory fragmentation is the production bottleneck:** naive pre-allocation wastes capacity on short requests; PagedAttention-style virtual memory is what makes long-context serving economically viable.
- **The trade-off is domain-specific:** KV caching only helps autoregressive decoders — apply it where the structural invariance holds, not as a generic optimization.

**Coming next:** Module 19 builds the statistical machinery that tells you whether the  $10\text{--}15\times$  speedup you just claimed is real — or just noise wearing a confidence interval.

## 28.13 Further Reading

For students who want to understand the academic foundations and production implementation of memoization in transformers:

### 28.13.1 Seminal Papers

- **Attention Is All You Need** - Vaswani et al. (2017). The original transformer paper that introduced the architecture requiring KV caching for efficient generation. Section 3.2 describes the attention mechanism that benefits from memoization. **Systems Implication:** The parallelization advantage over RNNs breaks the sequential compute bottleneck during training, but causes autoregressive inference to be heavily memory-bandwidth bound without caching. [arXiv:1706.03762](https://arxiv.org/abs/1706.03762)
- **Generating Sequences With Recurrent Neural Networks** - Graves (2013). Early work on autoregressive generation patterns, establishing the sequential token generation that creates the redundant computation KV caching eliminates. **Systems Implication:** RNN inference natively maintains a hidden state vector in memory ( $O(1)$  memory), making it highly efficient for single-batch generation on memory-constrained hardware. [arXiv:1308.0850](https://arxiv.org/abs/1308.0850)
- **Training Compute-Optimal Large Language Models** - Hoffmann et al. (2022). Analyzes the computational costs of training and inference, quantifying the importance of inference optimizations like KV caching at scale. **Systems Implication:** Established that optimal scaling requires massive datasets, turning training into a massive I/O problem where storage bandwidth and network interconnects dictate overall cluster efficiency. [arXiv:2203.15556](https://arxiv.org/abs/2203.15556)
- **FlashAttention: Fast and Memory-Efficient Exact Attention** - Dao et al. (2022). Modern attention optimization that combines with KV caching in production systems, demonstrating complementary optimization strategies. **Systems Implication:** Overcame GPU HBM bandwidth limits via SRAM tiling, fusing the attention computation into a single kernel to avoid costly memory reads and writes. [arXiv:2205.14135](https://arxiv.org/abs/2205.14135)

### 28.13.2 Additional Resources

- **System:** [vLLM documentation](#) - Production serving system that uses advanced KV cache management with paging
- **Tutorial:** [Hugging Face Text Generation Guide](#) - See `use_cache` parameter in production API
- **Blog:** [“The Illustrated Transformer”](#) by Jay Alammar - Visual explanation of attention mechanisms that benefit from caching

## 28.14 What’s Next

You’ve claimed a 10–15x speedup. The honest question is: how do you know? “Generation feels faster” is not an answer a systems engineer can defend. To put a real number on what you just built — and to compare it against the other optimizations in this tier — you need disciplined measurement.

### Coming Up: Module 19 - Benchmarking

Module 19 builds the measurement infrastructure that lets you say “this optimization gave us 12.4x speedup with 95% confidence” instead of “it seems faster”. You’ll implement statistical timing, warm-up handling, and Pareto-frontier analysis — the same tools production teams use to validate every optimization in this book, including the KV cache you just wrote.

**How memoization combines with the other optimizations in this tier:**

Table 28.6 traces how this module is reused by later parts of the curriculum.

Table 28.6: **How memoization stacks with quantization, acceleration, and benchmarking.**

| Module                  | What It Does                    | Works with Memoization                              |
|-------------------------|---------------------------------|-----------------------------------------------------|
| <b>15: Quantization</b> | Reduce precision to save memory | KVCache with int8 keys/values → 4x memory reduction |
| <b>17: Acceleration</b> | Optimize computation kernels    | Fused attention + KV cache → minimal memory traffic |
| <b>19: Benchmarking</b> | Measure end-to-end performance  | Profile cache hit rates and speedup gains           |

## 28.15 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## 🔥 Chapter 29

# Module 19: Benchmarking

A benchmark is not a number; it is a claim about throughput, latency, and hardware utilization under a specific workload. This module builds the measurement discipline that separates “my model is 2x faster” from “my model is 2x faster on batch=1, input-length=128, FP16, on an A100, with warmup discarded and 95% confidence intervals reported.” Every performance claim in an ML paper or an MLPerf submission lives or dies in the harness you write here.

### Module Info

**OPTIMIZATION TIER** | Difficulty: ●●●○ | Time: 5-7 hours | Prerequisites: 01-18

This module assumes familiarity with the complete TinyTorch stack (Modules 01-13), profiling (Module 14), and optimization techniques (Modules 15-18). You should understand how to build, profile, and optimize models before tackling systematic benchmarking and statistical comparison of optimizations.

## 29.1 Overview

“My model is 3x faster!” Faster than what? Measured how? At what input size, on what hardware, after how many warmup runs? Most “speedup” claims dissolve under those four questions — and yours will too if you don’t measure with care.

Modules 14-18 gave you optimizations. This module gives you the one tool that tells you which of them actually worked: a benchmarking harness that controls for noise, runs proper warmup, and produces confidence intervals you can defend. It is the same statistical machinery MLPerf uses, written in a few hundred lines you build yourself.

By the end, you will have the evaluation framework that drives the TorchPerf Olympics capstone — and the discipline to never publish a number you cannot justify.

## 29.2 Learning Objectives

### By completing this module, you will:

- **Implement** a benchmarking harness with confidence intervals, warmup protocols, and variance control
- **Quantify** trade-offs between accuracy, latency, and memory using Pareto frontiers
- **Diagnose** measurement noise — coefficient of variation, outliers, cold-start effects — before reporting numbers
- **Connect** optimizations from Modules 14-18 into a single comparison workflow for the TorchPerf Olympics capstone

## 29.3 What You'll Build




Figure 29.1: **TinyTorch Benchmarking System:** Statistical measurement of performance and accuracy.

### Implementation roadmap:

Table 29.1 lays out the implementation in order, one part at a time.

Table 29.1: **Implementation roadmap for the benchmarking suite.**

| Part | What You'll Implement                                 | Key Concept                                            |
|------|-------------------------------------------------------|--------------------------------------------------------|
| 1    | <code>BenchmarkResult</code> dataclass                | Statistical analysis with confidence intervals         |
| 2    | <code>precise_timer()</code> context manager          | High-precision timing for microsecond measurements     |
| 3    | <code>Benchmark.run_latency_benchmark()</code>        | Warmup protocols and latency measurement               |
| 4    | <code>Benchmark.run_accuracy_benchmark()</code>       | Model quality evaluation across datasets               |
| 5    | <code>Benchmark.run_memory_benchmark()</code>         | Peak memory tracking during inference                  |
| 6    | <code>BenchmarkSuite</code> for multi-metric analysis | Pareto frontier generation and trade-off visualization |

### The pattern you'll enable:

```

# Compare baseline vs optimized model with statistical rigor
benchmark = Benchmark([baseline_model, optimized_model])
latency_results = benchmark.run_latency_benchmark()
# Output: baseline: 12.3ms ± 0.8ms, optimized: 4.1ms ± 0.3ms (67% reduction, p <
0.01)
  
```

### 29.3.1 What You're NOT Building (Yet)

To keep this module focused, you will **not** implement:

- Hardware-specific benchmarks (GPU profiling requires CUDA, covered in production frameworks)
- Energy measurement (requires specialized hardware like power meters)
- Distributed benchmarking (multi-node coordination is beyond scope)
- Automated hyperparameter tuning for optimization (that is Module 20: Capstone)

**You are building the statistical foundation for fair comparison.** The hardware-specific extensions are mechanical once the methodology is right.

## 29.4 API Reference

This section documents the benchmarking API you will implement. Use it as your reference while building the statistical analysis and measurement infrastructure.

### 29.4.1 BenchmarkResult Dataclass

```
BenchmarkResult(metric_name: str, values: List[float], metadata: Dict[str, Any] = {})
```

Statistical container for benchmark measurements with automatic computation of mean, standard deviation, median, and 95% confidence intervals.

**Properties computed in `__post_init__`:**

The properties you expose are summarised in Table 29.2.

Table 29.2: Statistical properties computed by `BenchmarkResult`.

| Property              | Type               | Description                                  |
|-----------------------|--------------------|----------------------------------------------|
| <code>mean</code>     | <code>float</code> | Average of all measurements                  |
| <code>std</code>      | <code>float</code> | Standard deviation (0 if single measurement) |
| <code>median</code>   | <code>float</code> | Middle value, less sensitive to outliers     |
| <code>min_val</code>  | <code>float</code> | Minimum observed value                       |
| <code>max_val</code>  | <code>float</code> | Maximum observed value                       |
| <code>count</code>    | <code>int</code>   | Number of measurements                       |
| <code>ci_lower</code> | <code>float</code> | Lower bound of 95% confidence interval       |
| <code>ci_upper</code> | <code>float</code> | Upper bound of 95% confidence interval       |

**Methods:**

Table 29.3 lists the methods on the result dataclass.

Table 29.3: Serialization methods on the `BenchmarkResult` dataclass.

| Method               | Signature                                   | Description                             |
|----------------------|---------------------------------------------|-----------------------------------------|
| <code>to_dict</code> | <code>to_dict() -&gt; Dict[str, Any]</code> | Serialize to dictionary for JSON export |
| <code>__str__</code> | Returns formatted summary                   | "metric: mean ± std (n=count) "         |

### 29.4.2 Timing Context Manager

```
with precise_timer() as timer:
    # Your code to measure
    ...
# Access elapsed time
elapsed_seconds = timer.elapsed
```

High-precision timing context manager using `time.perf_counter()` for monotonic, nanosecond-resolution measurements.

### 29.4.3 Benchmark Class

```
Benchmark(models: List[Any], datasets: List[Any],
           warmup_runs: int = 5, measurement_runs: int = 10)
```

Core benchmarking engine for single-metric evaluation across multiple models.

**Parameters:** - `models`: List of models to benchmark (supports any object with forward/predict/call)  
 - `datasets`: List of datasets for accuracy benchmarking (required) - `warmup_runs`: Number of warmup iterations before measurement (default: 5) - `measurement_runs`: Number of measurement iterations for statistical analysis (default: 10)

**Core Methods:**

Table 29.4 lists the methods in this group.

Table 29.4: Core measurement methods on the `Benchmark` class.

| Method                              | Signature                                                                                         | Description                              |
|-------------------------------------|---------------------------------------------------------------------------------------------------|------------------------------------------|
| <code>run_latency_benchmark</code>  | <code>run_latency_benchmark(input_shape: Tuple[int, ...]) -&gt; Dict[str, BenchmarkResult]</code> | Measure inference time per model         |
| <code>run_accuracy_benchmark</code> | <code>run_accuracy_benchmark() -&gt; Dict[str, BenchmarkResult]</code>                            | Measure prediction accuracy on datasets  |
| <code>run_memory_benchmark</code>   | <code>run_memory_benchmark(input_shape: Tuple[int, ...]) -&gt; Dict[str, BenchmarkResult]</code>  | Track peak memory usage during inference |
| <code>compare_models</code>         | <code>compare_models(metric: str = "latency") -&gt; List[Dict]</code>                             | Compare models across a specific metric  |

### 29.4.4 BenchmarkSuite Class

```
BenchmarkSuite(models: List[Any], datasets: List[Any], output_dir: str =
               "benchmark_results")
```

Comprehensive multi-metric evaluation suite for generating Pareto frontiers and optimization trade-off analysis.

Table 29.5 lists the methods on the benchmark suite.

Table 29.5: Methods on the `BenchmarkSuite` class for multi-metric evaluation.

| Method                            | Signature                                                                     | Description                                        |
|-----------------------------------|-------------------------------------------------------------------------------|----------------------------------------------------|
| <code>run_full_benchmark</code>   | <code>run_full_benchmark() -&gt; Dict[str, Dict[str, BenchmarkResult]]</code> | Run all benchmark types and aggregate results      |
| <code>plot_pareto_frontier</code> | <code>plot_pareto_frontier(x_metric='latency', y_metric='accuracy')</code>    | Plot Pareto frontier for two competing objectives  |
| <code>plot_results</code>         | <code>plot_results(save_plots=True)</code>                                    | Generate visualization plots for benchmark results |
| <code>generate_report</code>      | <code>generate_report() -&gt; str</code>                                      | Generate comprehensive benchmark report            |

## 29.5 Core Concepts

This section covers the fundamental principles that make benchmarking scientifically valid. Understanding these concepts separates professional performance engineering from naive timing measurements.

### 29.5.1 Benchmarking Methodology

A single measurement tells you nothing. Time one inference and you might see 1.2ms; time it again and you might see 3.1ms because a background process woke up. Neither number is the model’s performance — they are samples from a noisy distribution.

Your job is to estimate that distribution’s mean with quantified confidence. That requires controlling the inputs you can (warmup, input shape, iteration count) and reporting the noise you cannot.

The methodology follows a structured protocol:

**Warmup Phase:** Modern systems adapt to workloads. JIT compilers optimize hot code paths after several executions. CPU frequency scales up under sustained load. Caches fill with frequently accessed data. Without warmup, your first measurements capture cold-start behavior, not steady-state performance.

```
# Warmup protocol in action
for _ in range(warmup_runs):
    _ = model(input_data)  # Run but discard measurements
```

**Measurement Phase:** After warmup, run the operation multiple times and collect timing data. The Central Limit Theorem tells us that with enough samples, the sample mean approaches the true mean, and we can compute confidence intervals.

Here is how the `Benchmark` class implements the complete protocol:

The code in `lst-19-benchmarking-latency-loop` makes this concrete.

```
def run_latency_benchmark(self, input_shape: Tuple[int, ...] = (1, 28, 28)) ->
    Dict[str, BenchmarkResult]:
    """Benchmark model inference latency using Profiler."""
    results = {}

    for i, model in enumerate(self.models):
        model_name = getattr(model, 'name', f'model_{i}')
```

```

latencies = []

# Create input tensor for this benchmark
input_data = Tensor(np.random.randn(*input_shape).astype(np.float32))

# Warmup runs (discard results)
for _ in range(self.warmup_runs):
    if hasattr(model, 'forward'):
        _ = model.forward(input_data)
    elif callable(model):
        _ = model(input_data)

# Measurement runs (collect statistics)
for _ in range(self.measurement_runs):
    with precise_timer() as timer:
        if hasattr(model, 'forward'):
            _ = model.forward(input_data)
        elif callable(model):
            _ = model(input_data)
    latencies.append(timer.elapsed)

results[model_name] = BenchmarkResult(
    metric_name=f"{model_name}_latency",
    values=latencies,
    metadata={'input_shape': input_shape, **self.system_info}
)

return results

```

: Listing 19.1 — Latency benchmark with warmup and measurement phases separated. {#lst-19-benchmarking-latency-loop}

Notice the clear separation between warmup (discarded) and measurement (collected) phases. This ensures measurements reflect steady-state performance.

**Statistical Analysis:** Raw measurements get transformed into confidence intervals that quantify uncertainty. The 95% confidence interval tells you: “If we ran this benchmark 100 times, 95 of those runs would produce a mean within this range.”

### 29.5.2 Metrics Selection

Different optimization goals require different metrics. Choosing the wrong metric leads to optimizing for the wrong objective.

**Latency (Time per Inference):** Measures how fast a model processes a single input. Critical for real-time systems like autonomous vehicles where a prediction must complete in 30ms before the next camera frame arrives. Measured in milliseconds or microseconds.

**Throughput (Inputs per Second):** Measures total processing capacity. Critical for batch processing systems like translating millions of documents. Higher throughput means more efficient hardware utilization. Measured in samples per second or frames per second. While throughput and latency both measure performance, maximizing one often severely degrades the other.

The tension between them is formalized by **Little’s Law**: for any stable queueing system,

$$L = \lambda W$$

where  $L$  is the average number of in-flight requests,  $\lambda$  is throughput (requests per second), and  $W$  is the average latency (seconds per request). The law is deceptively simple but fully general — it holds for any arrival process, any service-time distribution, and any scheduling policy. Its operational complexity is  $O(1)$ : you measure any two of the three, and the third follows. The consequence for ML serving is unavoidable: if you want to double throughput without doubling hardware, you must either halve latency (hard) or double the number of concurrent requests (which inflates latency unless the accelerator has spare parallelism). Benchmarking that reports only one of the three tells you nothing about where you actually sit on this curve.

#### **i Systems Implication: Throughput vs. Latency and Little's Law**

It is a common misconception that optimizing for latency automatically improves throughput. In modern hardware accelerators, these two metrics are fundamentally in tension. Processing inputs one-by-one minimizes per-request latency but leaves massive parallel vector units (ALUs) starved for data, resulting in abysmal GPU utilization. Conversely, batching inputs together drastically increases throughput and hardware utilization, but it artificially inflates latency because the first input must wait for the entire batch to finish processing. This trade-off is formalized in queueing theory by **Little's Law** ( $L = \lambda W$ ): to increase throughput ( $\lambda$ ) while keeping latency ( $W$ ) fixed, you must increase the number of requests processed concurrently ( $L$ ). Professional benchmarking is essential precisely because it maps where your system currently sits on this unforgiving mathematical curve.

**Accuracy (Prediction Quality):** Measures how often the model makes correct predictions. The fundamental quality metric. No point having a 1ms model if it is wrong 50% of the time. Measured as percentage of correct predictions on a held-out test set.

**Memory Footprint:** Measures peak RAM usage during inference. Critical for edge devices with limited memory. A 100MB model cannot run on a device with 64MB RAM. Measured in megabytes.

**Model Size:** Measures storage size of model parameters. Critical for over-the-air updates and storage-constrained devices. A 500MB model takes minutes to download on slow networks. Measured in megabytes.

The key insight: **these metrics trade off against each other**. Quantization reduces memory but may reduce accuracy. Pruning reduces latency but may require retraining. Professional benchmarking reveals these trade-offs quantitatively.

### 29.5.3 Statistical Validity

Without statistics, you cannot tell a real performance gap from noise — and shipping a “speedup” that is actually noise is how performance work loses credibility.

**Why Variance Matters:** Consider two models. Model A runs in 10.2ms, 10.3ms, 10.1ms (low variance). Model B runs in 9.5ms, 12.1ms, 8.9ms (high variance). Their means are identical, but B is unpredictable — a deal-breaker for any latency SLO. The mean alone hides this; the standard deviation surfaces it.

The `BenchmarkResult` class computes statistical properties automatically:

The code in `?@lst-19-benchmarking-result-dataclass` makes this concrete.

```
@dataclass
class BenchmarkResult:
    metric_name: str
    values: List[float]
    metadata: Dict[str, Any] = field(default_factory=dict)
```

```

def __post_init__(self):
    """Compute statistics after initialization."""
    if not self.values:
        raise ValueError("BenchmarkResult requires at least one measurement.")

    self.mean = statistics.mean(self.values)
    self.std = statistics.stdev(self.values) if len(self.values) > 1 else 0.0
    self.median = statistics.median(self.values)
    self.min_val = min(self.values)
    self.max_val = max(self.values)
    self.count = len(self.values)

    # 95% confidence interval for the mean
    if len(self.values) > 1:
        t_score = 1.96 # Approximate for large samples
        margin_error = t_score * (self.std / np.sqrt(self.count))
        self.ci_lower = self.mean - margin_error
        self.ci_upper = self.mean + margin_error
    else:
        self.ci_lower = self.ci_upper = self.mean

```

: Listing 19.2 — BenchmarkResult dataclass computing mean, std, median, and 95% CI. {#lst-19-benchmarking-result-dataclass}

The confidence interval calculation uses the standard error of the mean ( $\text{std} / \sqrt{n}$ ) scaled by the t-score. For large samples, a t-score of 1.96 corresponds to 95% confidence. This means: “We are 95% confident the true mean latency lies between ci\_lower and ci\_upper.”

**Coefficient of Variation:** The ratio  $\text{std} / \text{mean}$  measures relative noise. A CV of 0.05 means standard deviation is 5% of the mean, indicating stable measurements. A CV of 0.30 means 30% relative noise, indicating unstable or noisy measurements that need more samples.

**Outlier Detection:** Extreme values can skew the mean. The median is robust to outliers. If mean and median differ significantly, investigate outliers. They might indicate thermal throttling, background processes, or measurement errors.

### 29.5.4 Reproducibility

A benchmark that only you can reproduce is a benchmark nobody trusts. The minimum bar: a colleague pulls your code, runs it on equivalent hardware, and lands within your confidence interval.

**System Metadata:** Recording system configuration ensures results are interpreted correctly. A benchmark on a 2020 laptop will differ from a 2024 server, not because the model changed, but because the hardware did.

The code in ?@lst-19-benchmarking-system-info makes this concrete.

```

def __init__(self, models: List[Any], datasets: List[Any] = None,
             warmup_runs: int = DEFAULT_WARMUP_RUNS,
             measurement_runs: int = DEFAULT_MEASUREMENT_RUNS):
    self.models = models
    self.datasets = datasets or []
    self.warmup_runs = warmup_runs

```



```

self.measurement_runs = measurement_runs

# Capture system information for reproducibility
self.system_info = {
    'platform': platform.platform(),
    'python_version': platform.python_version(),
    'cpu_count': os.cpu_count() or 1,
}

```

: Listing 19.3 — Benchmark constructor capturing platform metadata for reproducibility. {#lst-19-benchmarking-system-info}

This metadata gets embedded in every BenchmarkResult, allowing you to understand why a benchmark run produced specific numbers.

**Controlled Environment:** Background processes, thermal state, and power settings all affect measurements. Professional benchmarking controls these factors:

- Close unnecessary applications before benchmarking
- Let the system reach thermal equilibrium (run warmup)
- Use the same hardware configuration across runs
- Document any anomalies or environmental changes

**Versioning:** Record the versions of all dependencies. A NumPy update might change BLAS library behavior, affecting performance. Recording versions ensures results remain interpretable months later.

### 29.5.5 Reporting Results

A benchmark that nobody can act on is wasted compute. Reports must do two jobs: surface the trade-off, and be defensible under questioning.

**Comparison Tables:** Show mean, standard deviation, and confidence intervals for each model on each metric. This lets stakeholders quickly identify winners and understand uncertainty.

**Pareto Frontiers:** When metrics trade off, visualize the Pareto frontier to show which models are optimal for different constraints. A model is Pareto-optimal if no other model is better on all metrics simultaneously.

Consider three models: - Model A: 10ms latency, 90% accuracy - Model B: 15ms latency, 95% accuracy - Model C: 12ms latency, 91% accuracy

Model C is dominated by Model A (faster and only 1% less accurate). It is not Pareto-optimal. Models A and B are both Pareto-optimal: if you want maximum accuracy, choose B; if you want minimum latency, choose A.

**Visualization:** Scatter plots reveal relationships between metrics. Plot latency vs accuracy and you immediately see the trade-off frontier. Add annotations showing model names and you have an actionable decision tool.

**Statistical Significance:** Report not just means but confidence intervals. Saying “Model A is 2ms faster” is incomplete. Saying “Model A is 2ms faster with 95% confidence interval [1.5ms, 2.5ms],  $p < 0.01$ ” provides the statistical rigor needed for engineering decisions.

## 29.6 Common Errors

Four mistakes account for almost every misleading benchmark you will encounter — in classrooms, in blog posts, and in production. Spot them in your own code first.

### 29.6.1 Insufficient Measurement Runs

**Error:** Running a benchmark only 3 times produces unreliable statistics.

**Symptom:** Results change dramatically between benchmark runs. Confidence intervals are extremely wide.

**Cause:** The Central Limit Theorem requires sufficient samples. With only 3 measurements, the sample mean is a poor estimator of the true mean.

**Fix:** Use at least 10 measurement runs (the default in this module). For high-variance operations, increase to 30+ runs.

```
# BAD: Only 3 measurements
benchmark = Benchmark(models, measurement_runs=3) # Unreliable!

# GOOD: 10+ measurements
benchmark = Benchmark(models, measurement_runs=10) # Statistical confidence
```

### 29.6.2 Skipping Warmup

**Error:** Measuring performance without warmup captures cold-start behavior, not steady-state performance.

**Symptom:** First measurement is much slower than subsequent ones. Results do not reflect production performance.

**Cause:** JIT compilation, cache warming, and CPU frequency scaling all require several iterations to stabilize.

**Fix:** Always run warmup iterations before measurement.

```
# Already handled in Benchmark class
for _ in range(self.warmup_runs):
    _ = model(input_data) # Warmup (discarded)

for _ in range(self.measurement_runs):
    # Now measure steady-state performance
```

### 29.6.3 Comparing Different Input Shapes

**Error:** Benchmarking Model A on 28x28 images and Model B on 224x224 images, then comparing latency.

**Symptom:** Misleading conclusions about which model is faster.

**Cause:** Larger inputs require more computation. You are measuring input size effects, not model efficiency.

**Fix:** Use identical input shapes across all models in a benchmark.

```
# Ensure all models use same input shape
results = benchmark.run_latency_benchmark(input_shape=(1, 28, 28))
```

### 29.6.4 Ignoring Variance

**Error:** Reporting only the mean, ignoring standard deviation and confidence intervals.

**Symptom:** Cannot determine if performance differences are statistically significant or just noise.

**Cause:** Treating measurements as deterministic when they are actually stochastic.

**Fix:** Always report confidence intervals, not just means.

```
# BenchmarkResult automatically computes confidence intervals
result = BenchmarkResult("latency", measurements)
print(f"{result.mean:.3f}ms ± {result.std:.3f}ms, 95% CI: [{result.ci_lower:.3f},
      {result.ci_upper:.3f}]" )
```

## 29.7 Production Context

### 29.7.1 Your Implementation vs. Industry Benchmarks

Your TinyTorch benchmarking infrastructure implements the same statistical principles used in production ML benchmarking frameworks. The difference is scale and automation.

Table 29.6 places your implementation side by side with the production reference for direct comparison.

Table 29.6: Feature comparison between TinyTorch benchmarking and MLPerf practices.

| Feature                     | Your Implementation       | MLPerf / Industry                        |
|-----------------------------|---------------------------|------------------------------------------|
| <b>Statistical Analysis</b> | Mean, std, 95% CI         | Same + hypothesis testing, ANOVA         |
| <b>Metrics</b>              | Latency, accuracy, memory | Same + energy, throughput, tail latency  |
| <b>Warmup Protocol</b>      | Fixed warmup runs         | Same + adaptive warmup until convergence |
| <b>Reproducibility</b>      | System metadata           | Same + hardware specs, thermal state     |
| <b>Automation</b>           | Manual benchmark runs     | CI/CD integration, regression detection  |
| <b>Scale</b>                | Single machine            | Distributed benchmarks across clusters   |

### 29.7.2 Code Comparison

The following shows equivalent benchmarking patterns in TinyTorch and production frameworks like MLPerf.

## 29.8 Your TinyTorch

```
from tinytorch.perf.benchmarking import Benchmark

# Setup models and benchmark
benchmark = Benchmark(
    models=[baseline_model, optimized_model],
    warmup_runs=5,
    measurement_runs=10
)

# Run latency benchmark
results = benchmark.run_latency_benchmark(input_shape=(1, 28, 28))
```

```
# Analyze results
for model_name, result in results.items():
    print(f"{model_name}: {result.mean*1000:.2f}ms ± {result.std*1000:.2f}ms")
```

## 29.9 MLPerf (Industry Standard)

```
import mlperf_loadgen as lg

# Configure benchmark scenario
settings = lg.TestSettings()
settings.scenario = lg.TestScenario.SingleStream
settings.mode = lg.TestMode.PerformanceOnly

# Run standardized benchmark
sut = SystemUnderTest(model)
lg.StartTest(sut, qsl, settings)

# Results include latency percentiles, throughput, accuracy
```

Let's understand the comparison:

- **Line 1-3 (Setup):** TinyTorch uses a simple class-based API. MLPerf uses a loadgen library with standardized scenarios (SingleStream, Server, MultiStream, Offline). Both ensure fair comparison.
- **Line 6-8 (Configuration):** TinyTorch exposes `warmup_runs` and `measurement_runs` directly. MLPerf abstracts this into `TestSettings` with scenario-specific defaults. Same concept, different abstraction level.
- **Line 11-13 (Execution):** TinyTorch returns `BenchmarkResult` objects with statistics. MLPerf logs results to standardized formats that compare across hardware vendors. Both provide statistical analysis.
- **Statistical Rigor:** Both use repeated measurements, warmup, and confidence intervals. TinyTorch teaches the foundations; MLPerf adds industry-specific requirements.

### 💡 What's Identical

The statistical methodology, warmup protocols, and reproducibility requirements are identical. Production frameworks add automation, standardization across organizations, and hardware-specific optimizations. Understanding TinyTorch benchmarking gives you the foundation to work with any industry benchmarking framework.

### 29.9.1 Why Benchmarking Matters at Scale

At production scale, small performance differences compound into massive resource consumption — which is exactly why honest measurement is non-negotiable:

- **Cost:** A data center running 10,000 GPUs 24/7 consumes roughly \$50 million in electricity annually. A real 10% latency reduction saves \$5 million per year. A *fake* 10% — one that does not survive a careful benchmark — wastes the engineering budget that chased it.
- **User Experience:** Search engines must return results in under 200ms. A 50ms latency reduction is the difference between keeping or losing users.

- **Sustainability:** Training GPT-3 consumed 1,287 MWh, equivalent to the annual energy use of 120 US homes. Optimization reduces carbon footprint — but only if the “optimization” is real.

This is why fair benchmarking is a discipline, not a courtesy. The next chapter — the capstone — puts that discipline to the test on a public leaderboard.

## 29.10 Check Your Understanding

### 💡 Check Your Understanding — Benchmarking

Before moving on, verify you can articulate each of the following:

- Little’s Law ( $L = \lambda W$ ) applied to throughput vs latency trade-offs under fixed parallelism, and why you cannot improve both without adding concurrency.
- Why a single measurement is worthless: warmup, variance, and 95% confidence intervals turn a sample into a defensible claim.
- How to read a Pareto frontier across (latency, accuracy, memory) and identify which models are dominated vs optimal.
- The four measurement traps (wrong thing, system noise, cold start, batch-size confusion) and the specific protocol that defuses each.

If any of these feels fuzzy, revisit the Core Concepts section (especially Benchmarking Methodology and Statistical Validity) before moving on.

Test your understanding of benchmarking statistics and methodology with these quantitative questions. Work each one with a calculator before opening the answer.

#### Q1: Statistical Significance

You benchmark a baseline model and an optimized model 10 times each. Baseline: mean=12.5ms, std=1.2ms. Optimized: mean=11.8ms, std=1.5ms. Is the optimized model statistically significantly faster?

### 💡 Answer

**Calculate 95% confidence intervals:**

Baseline:  $CI = \text{mean} \pm 1.96 * (\text{std} / \text{sqrt}(n)) = 12.5 \pm 1.96 * (1.2 / \text{sqrt}(10)) = 12.5 \pm 0.74 = [11.76, 13.24]$

Optimized:  $CI = 11.8 \pm 1.96 * (1.5 / \text{sqrt}(10)) = 11.8 \pm 0.93 = [10.87, 12.73]$

**Result:** The confidence intervals OVERLAP (baseline goes as low as 11.76, optimized goes as high as 12.73). The difference is **NOT statistically significant** at the 95% confidence level. You cannot claim the optimized model is faster.

**Lesson:** Always compute confidence intervals. A 0.7ms difference in means looks meaningful, but with these variances and sample sizes it could be random noise.

#### Q2: Sample Size Calculation

You measure latency with standard deviation of 2.0ms. How many measurements do you need to achieve a 95% confidence interval width of  $\pm 0.5$ ms?

### 💡 Answer

**Confidence interval formula:**  $\text{margin} = 1.96 * (\text{std} / \text{sqrt}(n))$

**Solve for n:**  $0.5 = 1.96 * (2.0 / \text{sqrt}(n))$

$\text{sqrt}(n) = 1.96 * 2.0 / 0.5 = 7.84$

$n = 7.84^2 = 61.5 \approx 62$  measurements

**Lesson:** Achieving tight confidence intervals requires many measurements. Quadrupling precision (from  $\pm 1.0\text{ms}$  to  $\pm 0.5\text{ms}$ ) requires 4x more samples (15 to 60). This is why professional benchmarks run hundreds of iterations.

### Q3: Warmup Impact

Without warmup, your measurements are: [15.2, 12.1, 10.8, 10.5, 10.6, 10.4] ms. With 3 warmup runs discarded, your measurements are: [10.5, 10.6, 10.4, 10.7, 10.5, 10.6] ms. How much does warmup reduce measured latency and variance?

#### 💡 Answer

##### Without warmup:

- Mean =  $(15.2 + 12.1 + 10.8 + 10.5 + 10.6 + 10.4) / 6 = 11.6\text{ms}$
- Std = 1.8ms (high variance from warmup effects)

##### With warmup:

- Mean =  $(10.5 + 10.6 + 10.4 + 10.7 + 10.5 + 10.6) / 6 = 10.55\text{ms}$
- Std = 0.1ms (low variance, stable measurements)

##### Impact:

- Latency reduced:  $11.6\text{ms} - 10.55\text{ms} = 1.05\text{ms}$  (9% reduction)
- Variance reduced:  $1.8 \rightarrow 0.1\text{ms} = 94\%$  reduction in noise

**Lesson:** Warmup eliminates cold-start effects and dramatically reduces measurement variance. Without warmup, you are measuring system startup behavior, not steady-state performance.

### Q4: Pareto Frontier

You have three models with (latency, accuracy): A=(5ms, 88%), B=(8ms, 92%), C=(6ms, 89%). Which are Pareto-optimal?

#### 💡 Answer

**Pareto-optimal definition:** A model is Pareto-optimal if no other model is better on ALL metrics simultaneously.

**Analysis:** - Model A vs B: A is faster ( $5\text{ms} < 8\text{ms}$ ) but less accurate ( $88\% < 92\%$ ). Neither dominates. Both Pareto-optimal. - Model A vs C: A is faster ( $5\text{ms} < 6\text{ms}$ ) but less accurate ( $88\% < 89\%$ ). Neither dominates. Both Pareto-optimal. - Model B vs C: B is slower ( $8\text{ms} > 6\text{ms}$ ) but more accurate ( $92\% > 89\%$ ). Neither dominates. Both Pareto-optimal.

**Result:** All three models are Pareto-optimal. None is strictly dominated by another.

**Lesson:** The Pareto frontier shows trade-off options. If you need minimum latency, choose A. If you need maximum accuracy, choose B. If you want balanced performance, choose C.

### Q5: Measurement Overhead

Your timer has  $1\mu\text{s}$  overhead per measurement. You measure a  $50\mu\text{s}$  operation 1000 times. What percentage of measured time is overhead?

#### 💡 Answer

**Total true operation time:**  $50\mu\text{s} \times 1000 = 50,000\mu\text{s} = 50\text{ms}$

**Total timer overhead:**  $1\mu\text{s} \times 1000 = 1,000\mu\text{s} = 1\text{ms}$

**Total measured time:**  $50\text{ms} + 1\text{ms} = 51\text{ms}$

**Overhead percentage:**  $(1\text{ms} / 51\text{ms}) \times 100\% = 1.96\%$

**Lesson:** Timer overhead is negligible for operations longer than  $\sim 50\mu\text{s}$  but dominates for microsecond-scale operations. That is why we use `time.perf_counter()` (nanosecond resolution, minimal overhead). For operations under  $10\mu\text{s}$ , time a batch and divide.

## 29.11 Key Takeaways

- **$L = \lambda W$  is the serving-curve law:** pick any two of (concurrency, throughput, latency) and the third follows — optimizing one in isolation is meaningless.
- **Variance is the signal, not the noise:** a “faster” mean inside an overlapping confidence interval is not faster; reporting std and 95% CI is the floor, not extra credit.
- **Warmup is mandatory:** JIT compilation, cache warming, and CPU frequency scaling all skew the first measurements — discard them, or you are benchmarking initialization, not steady state.
- **Pareto frontiers beat single-axis wins:** a  $2\times$  speedup that costs 5% accuracy may not be a win; only the frontier view reveals which optimizations dominate the others.

**Coming next:** Module 20 puts the harness to work — the TorchPerf Olympics capstone asks you to stack every optimization you built and defend the resulting numbers against the same statistical scrutiny.

## 29.12 Further Reading

For students who want to understand the academic and industry foundations of ML benchmarking:

### 29.12.1 Seminal Papers

- **MLPerf: An Industry Standard Benchmark Suite for Machine Learning Performance** - Mattson et al. (2020). The definitive industry benchmark framework that establishes fair comparison methodology across hardware vendors. Your benchmarking infrastructure follows the same statistical principles MLPerf uses for standardized evaluation. **Systems Implication:** Standardized rules around hardware-specific optimizations (like mixed precision and custom kernels), enforcing strict timing boundaries to measure true system-level end-to-end throughput. [arXiv:1910.01500](#)
- **How to Benchmark Code Execution Times on Intel IA-32 and IA-64 Instruction Set Architectures** - Paoloni (2010). White paper explaining low-level timing mechanisms, measurement overhead, and sources of timing variance. Essential reading for understanding what happens when you call `time.perf_counter()`. **Systems Implication:** Highlighted the impact of CPU clock scaling, thread migration, and caching effects, requiring developers to flush caches and pin threads to get stable nanosecond-level measurements. [Intel](#)
- **Statistical Tests for Comparing Performance** - Georges et al. (2007). Academic treatment of statistical methodology for benchmarking, including hypothesis testing, sample size calculation, and handling measurement noise. **Systems Implication:** Demonstrated that garbage collection pauses, JIT compilation, and OS interrupts introduce non-deterministic latency, necessitating multi-trial warmups to measure steady-state hardware performance. [ACM OOPSLA](#)
- **Benchmarking Deep Learning Inference on Mobile Devices** - Ignatov et al. (2018). Comprehensive study of mobile ML benchmarking challenges including thermal throttling, battery constraints, and heterogeneous hardware. Shows how benchmarking methodology changes for resource-constrained devices. **Systems Implication:** Mobile deployment is constrained by thermal dissipation and battery limits, meaning sustained compute power drops significantly over time, heavily favoring integer quantization over floating-point operations. [arXiv:1812.01328](#)

### 29.12.2 Additional Resources

- **Industry Standard:** [MLCommons MLPerf](#) - Browse actual benchmark results across hardware vendors to see professional benchmarking in practice
- **Textbook:** “The Art of Computer Systems Performance Analysis” by Raj Jain - Comprehensive treatment of experimental design, statistical analysis, and benchmarking methodology for systems engineering

## 29.13 What’s Next

You now have a tool that separates real speedups from noise. Before the Capstone, the next chapter puts the whole Optimization Tier through a historical end-to-end run. The **Optimization Milestone** — MLPerf (2018) — chains your Profiler (Module 14), Quantization (15), Compression (16), Acceleration (17), and KV-Cache (18) into a single measure → optimize → validate loop, the same discipline MLPerf forced onto an industry that previously measured whatever made its hardware look fastest. You compress a model 8× and speed up transformer generation 10×, on your own framework, with every speedup traceable to code you wrote.

After that, Module 20 — the **Capstone: TorchPerf Olympics** — is where the benchmarking harness earns its keep. You will combine the optimizations from Modules 14-18 (quantization, pruning, fusion, caching), submit results to a public leaderboard, and defend every number against the same statistical scrutiny you just built.

The capstone has no “easy” mode. Every entry is judged by the harness from this chapter — overlapping confidence intervals are not a win, and a faster mean without proper warmup is a disqualification. The discipline you practiced here is the price of admission.

#### **i** Coming Up: Optimization Milestone (MLPerf), then Module 20 — Capstone

First: the MLPerf milestone runs your Profiler → Quantization → Pruning → KV-Cache pipeline on real models, the same measure-optimize-validate loop production teams use. Then Module 20 combines everything from Modules 01-19 to compete in the TorchPerf Olympics: stack optimizations, benchmark them honestly, and chase the fastest, smallest, or most accurate model on the leaderboard.

#### **Preview — How Your Benchmarking Gets Used in the Capstone:**

Table 29.7 traces how this module is reused by later parts of the curriculum.

Table 29.7: **How benchmarking powers each capstone competition event.**

| Competition Event       | Metric Optimized                    | Your Benchmark In Action                                              |
|-------------------------|-------------------------------------|-----------------------------------------------------------------------|
| <b>Latency Sprint</b>   | Minimize inference time             | <code>benchmark.run_latency_benchmark()</code><br>determines winners  |
| <b>Memory Challenge</b> | Minimize model size                 | <code>benchmark.run_memory_benchmark()</code><br>tracks footprint     |
| <b>Accuracy Contest</b> | Maximize accuracy under constraints | <code>benchmark.run_accuracy_benchmark()</code><br>validates quality  |
| <b>All-Around</b>       | Balanced Pareto frontier            | <code>suite.plot_pareto_frontier()</code> finds<br>optimal trade-offs |



## 29.14 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



---

PART VI

# Optimization Milestones

---



## 🔥 Chapter 30

# Milestone 06: MLPerf — The Optimization Era (2018)

## i Milestone Info

**Optimization Milestone** | Difficulty: ●●●● | Time: 1–2 hours | Prerequisites: Modules 01–19

## 💡 What You'll Learn

- The systematic optimization workflow: measure, optimize, validate, repeat
- Why profiling before optimizing beats heroic rewrites
- How to achieve 8× compression and 10× speedup with minimal accuracy loss

## 30.1 Overview

This is the **Optimization Milestone** — the third and final act of the historical arc. The Foundation Milestones proved your training loop learns. The Architecture Milestones proved your layers match real data. This one proves your *optimization stack* — the Profiler (Module 14), Quantization (15), Compression (16), Acceleration (17), and KV-Cache (18) you just finished — can turn a working model into a shippable one.

2018. ML research is sprinting; deployment is crawling. BERT-Large lands at 340M parameters and won't fit on a phone. ResNet-152 inference blows past production latency budgets. Teams ship two-year-old models because the new ones are too slow, too big, too expensive — and every vendor's benchmark numbers use a different dataset, batch size, and accuracy floor, so you can't even tell whose hardware is faster.

MLPerf changes that. One protocol, one accuracy floor, one set of reference models — apples-to-apples across CPUs, GPUs, TPUs, and accelerators that didn't exist last year. Optimization stops being an afterthought and becomes the discipline that decides who ships.

You compress YOUR models 8× and accelerate YOUR transformer generation 10× using the same measure → optimize → validate loop production teams run. That's the gap between a research demo and a shipped product, closed in two scripts.

## 30.2 What You'll Build

A complete MLPerf-style optimization pipeline:

1. **Static model optimization** — profile, quantize, and prune MLP/CNN
2. **Generation speedup** — KV-cache acceleration for transformers

Measure --> Optimize --> Validate --> Repeat

### 30.3 Prerequisites

Table 30.1 lists the modules you need to have completed before starting.

Table 30.1: Prerequisite modules for the MLPerf milestone.

| Module | Component              | What It Provides                              |
|--------|------------------------|-----------------------------------------------|
| 01–08  | Foundation + Training  | Models and training loop to optimize          |
| 11–12  | Embeddings + Attention | Transformer components for generation speedup |
| 14     | Profiling              | YOUR measurement tools                        |
| 15     | Quantization           | YOUR INT8/FP16 implementations                |
| 16     | Compression            | YOUR pruning techniques                       |
| 17     | Acceleration           | YOUR vectorized operations                    |
| 18     | Memoization            | YOUR KV-cache for generation                  |
| 19     | Benchmarking           | YOUR standardized benchmark reports           |

### 30.4 Running the Milestone

Before running, ensure you have completed Modules 01–19. You can check your progress:

```
tito module status
```

```
cd milestones/06_2018_mlperf

# Part 1: Optimize MLP/CNN (profiling + quantization + pruning)
python 01_optimization_olympics.py
# Expected: 4–8x compression with <2% accuracy loss

# Part 2: Speed up Transformer generation (KV caching)
python 02_generation_speedup.py
# Expected: 6–10x faster generation
```

### 30.5 Expected Results

#### Static model optimization (script 01)

Table 30.2 tracks model size and accuracy through each static optimisation stage.

Table 30.2: Expected size and accuracy at each static optimization stage.

| Optimization          | Size  | Accuracy | Notes            |
|-----------------------|-------|----------|------------------|
| Baseline (FP32)       | 100%  | 85–90%   | Full precision   |
| + Quantization (INT8) | 25%   | 84–89%   | 4× smaller       |
| + Pruning (50%)       | 12.5% | 82–87%   | 8× smaller total |

#### Generation speedup (script 02)

Table 30.3 shows per-token generation speed with and without the KV cache.

Table 30.3: Per-token generation speed with and without the KV cache.

| Mode             | Time/Token | Speedup |
|------------------|------------|---------|
| Without KV-Cache | ~10 ms     | 1×      |
| With KV-Cache    | ~1 ms      | 6–10×   |

## 30.6 The Aha Moment: Systematic Beats Heroic

### The wrong way (heroic optimization):

"It's too slow! Let me rewrite everything in C++!"

"Memory is too high! Let me redesign the architecture!"

"KV-cache sounds complex! Let me try CUDA kernels first!"

Result: weeks of work, marginal gains, introduced bugs.

### The right way (systematic optimization):

1. MEASURE: Profile shows 70% of time is in attention, 80% of memory is Linear layers
2. OPTIMIZE: Add KV-cache (targets the 70%), quantize Linear layers (targets the 80%)
3. VALIDATE: Accuracy drops 1.5% (acceptable), 8x faster (huge win)
4. REPEAT: Profile again, find next bottleneck

Result: 10× faster, 8× smaller, 2% accuracy cost — achieved in days.

### This is what separates ML researchers from ML engineers:

- YOUR Profiler (Module 14) identifies real bottlenecks (not assumed ones)
- YOUR Quantization (Module 15) reduces memory 4×
- YOUR Pruning (Module 16) reduces parameters 50%+
- YOUR KV-Cache (Module 18) speeds up generation 10×

The full loop — measure, optimize, validate — runs on YOUR tools, not someone else's library.

## 30.7 Your Code Powers This

This milestone closes the historical arc. Every optimization tool you exercise here comes from YOUR implementations:

Table 30.4 names the TinyTorch components that power this milestone.

Table 30.4: TinyTorch components that power the MLPerf milestone.

| Component                | Your Module | What It Does                                   |
|--------------------------|-------------|------------------------------------------------|
| Profiler                 | Module 14   | YOUR measurement and bottleneck identification |
| <code>quantize()</code>  | Module 15   | YOUR INT8/FP16 conversion                      |
| <code>prune()</code>     | Module 16   | YOUR weight pruning                            |
| <code>vectorize()</code> | Module 17   | YOUR accelerated operations                    |

| Component | Your Module | What It Does                          |
|-----------|-------------|---------------------------------------|
| KVCache   | Module 18   | YOUR key-value caching for generation |

No external optimization libraries. Every speedup, every byte saved, traces back to code you wrote.

## 30.8 Historical Context

Before MLPerf, comparing ML systems was guesswork. Vendors picked their own datasets, batch sizes, and accuracy targets, then claimed wins. MLPerf forced a common protocol — same models, same data, same accuracy floor — so a “ $2\times$  faster” claim could finally be checked instead of believed.

That protocol marks the moment ML engineering became as load-bearing as ML research. Building a model is step one. Shipping it inside a latency budget, on hardware your users actually own, is where production value lives — and where careers are made.

## 30.9 Systems Insights

- **Memory** —  $4\text{--}16\times$  compression with under 2% accuracy loss is routine, not heroic.
- **Latency** —  $10\text{--}40\times$  speedup from caching and batching alone, before touching the model.
- **Trade-offs** — every win costs accuracy, latency, or memory. Profiling tells you which one to spend; intuition will lie to you.

## 30.10 What’s Next

With Milestone 06 the historical arc closes. Six recreations, nineteen modules, one framework you wrote yourself end-to-end. You’ve:

- Built every core component (Modules 01–13)
- Optimized and benchmarked for production deployment (Modules 14–19)
- Proven mastery by recreating six landmark systems (Milestones 01–06)

What’s left is the test no recreation can give you. The **Capstone** (Module 20) is the Torch Olympics — open-ended problems, fixed budgets, your framework, your call. The history lessons end here. The competition starts next.

## 30.11 Further Reading

- **MLPerf:** [mlcommons.org](https://mlcommons.org)
- **Deep Compression:** Han et al. (2015). “Deep Compression: Compressing DNNs with Pruning, Trained Quantization and Huffman Coding”
- **Efficient Transformers:** Tay et al. (2020). “Efficient Transformers: A Survey”



---

**PART VII**

# Capstone

---



## 🔥 Chapter 31

# Module 20: Capstone

The framework you wrote now has to run end-to-end, on one model, with every optimization stacked and every number defended. The capstone is where profiling, quantization, compression, acceleration, memoization, and benchmarking land on a single inference pipeline and produce an MLPerf-style `results.json` with platform metadata, warmup-controlled latencies, and schema-validated confidence intervals. Optimizations interact at runtime; this is where you measure whether they actually compound on silicon.

### Module Info

**OPTIMIZATION TIER** | Difficulty: ●●●● | Time: 6-8 hours | Prerequisites: All modules (01-19)

This capstone assumes you've built a complete framework end-to-end:

- Core framework (Modules 01-13) — **required**
- Optimization techniques (Modules 14-18) — **recommended**
- Benchmarking methodology (Module 19) — **required**

Parts 1-4 run on Modules 01-13 + 19 alone. Part 4b layers in Modules 14-18 to drive the full baseline-to-optimized workflow; without them, the pipeline degrades to baseline benchmarking only.

## 31.1 Overview

Nineteen modules in, you have a working ML framework. Tensors, autograd, transformers, quantization, pruning, profiling — every line written by you. The Olympics is where you put it on the scale.

In production ML, an unmeasured claim is no claim at all. This capstone gives your framework the same treatment MLPerf and Papers with Code give a real submission: a baseline measurement, an optimization pass, an apples-to-apples comparison, and a schema-validated artifact someone else can verify. You'll write the benchmarking harness, run it against your own code, and ship a `results.json` that stands on its own.

When you finish, you won't just have a framework. You'll have evidence.

## 31.2 Learning Objectives

### By completing this module, you will:

- **Implement** comprehensive benchmarking infrastructure measuring accuracy, latency, throughput, and memory
- **Master** the three pillars of reliable benchmarking: repeatability, comparability, and completeness
- **Understand** performance measurement traps (variance, cold starts, batch effects) and how to avoid them
- **Connect** your TinyTorch implementation to production ML workflows (experiment tracking, A/B testing, regression detection)

- **Generate** schema-validated JSON submissions that enable reproducible comparisons and community sharing

### 31.3 What You'll Build




Figure 31.1: **TinyTorch Submission Pipeline:** Baseline and optimized models flow through `BenchmarkReport` and `generate_submission()` to produce a schema-validated `results.json`.

**Implementation roadmap:**

Table 31.1 lays out the implementation in order, one part at a time.

Table 31.1: **Implementation roadmap for the capstone benchmarking artifact.**

| Part | What You'll Implement                     | Key Concept                                                  |
|------|-------------------------------------------|--------------------------------------------------------------|
| 1    | <code>SimpleMLP</code>                    | Baseline model for benchmarking demonstrations               |
| 2    | <code>BenchmarkReport</code>              | Comprehensive performance measurement with statistical rigor |
| 3    | <code>generate_submission()</code>        | Standardized JSON generation with schema compliance          |
| 4    | <code>validate_submission_schema()</code> | Automated validation ensuring data quality                   |
| 5    | Complete workflows                        | Baseline, optimization, comparison, submission pipeline      |

**The pattern you'll enable:**

```
# Professional ML workflow
report = BenchmarkReport(model_name="my_model")
report.benchmark_model(model, X_test, y_test, num_runs=100)
```

```
submission = generate_submission(
    baseline_report=baseline_report,
    optimized_report=optimized_report,
    techniques_applied=["quantization", "pruning"]
)
save_submission(submission, "results.json")
```

### 31.3.1 What You're NOT Building (Yet)

Out of scope for this capstone:

- CI/CD pipelines that benchmark on every commit
- Multi-hardware comparison across CPU/GPU/TPU
- Plotting dashboards for accuracy-vs-latency curves
- Leaderboard aggregation across community submissions

You are building the measurement and reporting foundation everything else stands on. Automation and visualization layer on top — but only if the numbers underneath are trustworthy.

## 31.4 API Reference

This section provides a quick reference for the benchmarking classes and functions you'll build. Use this while implementing and debugging.

### 31.4.1 BenchmarkReport Constructor

```
BenchmarkReport(model_name: str = "model") -> BenchmarkReport
```

Creates a benchmark report instance that measures and stores model performance metrics along with system context for reproducibility.

### 31.4.2 BenchmarkReport Properties

The properties you expose are summarised in Table 31.2.

Table 31.2: Data properties stored on the **BenchmarkReport** class.

| Property    | Type | Description                                        |
|-------------|------|----------------------------------------------------|
| model_name  | str  | Identifier for the model being benchmarked         |
| metrics     | dict | Performance measurements (accuracy, latency, etc.) |
| system_info | dict | Platform, Python version, NumPy version            |
| timestamp   | str  | When benchmark was run (ISO format)                |

### 31.4.3 Core Methods

Table 31.3 lists the methods on the report class.

Table 31.3: Core methods on the BenchmarkReport class.

| Method                     | Signature                                                                                | Description                                                      |
|----------------------------|------------------------------------------------------------------------------------------|------------------------------------------------------------------|
| benchmark_model            | <code>benchmark_model(model, X_test, y_test, num_runs=100)</code><br>-> dict             | Measures accuracy, latency (mean $\pm$ std), throughput, memory  |
| generate_submission        | <code>generate_submission(baseline_report, optimized_report=None, ...)</code><br>-> dict | Creates standardized JSON with baseline, optimized, improvements |
| save_submission            | <code>save_submission(submission, filepath="submission.json")</code> -> str              | Writes JSON to file with validation                              |
| validate_submission_schema | <code>validate_submission_schema(submission)</code><br>-> bool                           | Validates structure and value ranges                             |

### 31.4.4 Module Dependencies and Imports

This capstone integrates components from across TinyTorch:

#### Core dependencies (required):

```
from tinytorch.core.tensor import Tensor
from tinytorch.core.layers import Linear
from tinytorch.core.activations import ReLU
from tinytorch.core.losses import CrossEntropyLoss
```

#### Optimization modules (optional):

```
# These imports use try/except blocks for graceful degradation
try:
    from tinytorch.perf.profiling import Profiler, quick_profile
    from tinytorch.perf.compression import magnitude_prune, structured_prune
    from tinytorch.perf.benchmarking import Benchmark, BenchmarkResult
except ImportError:
    # Core benchmarking still works without optimization modules
    pass
```

The advanced optimization workflow (Part 4b) demonstrates these optional integrations, but the core benchmarking system (Parts 1-4) works with just the foundation modules (01-13) and basic benchmarking (19).

## 31.5 Core Concepts

This section covers the fundamental principles of professional ML benchmarking. These concepts apply to every ML system, from research papers to production deployments.

### 31.5.1 The Reproducibility Crisis in ML

ML has a credibility problem. A paper claims “92% accuracy with 10ms latency” — and you can’t reproduce it, because the paper never said which hardware, which software versions, which batch size, or how the measurement was taken. Without those facts the number is folklore.

MLPerf and Papers with Code emerged to fix this, and they did it by requiring four things from every submission:

- **Standardized tasks** on fixed datasets
- **Hardware specifications** documented in full
- **Measurement protocols** pinned down precisely
- **Code submissions** that anyone can re-run

Your benchmarking system enforces the same contract. Every submission you generate carries the context another person needs to reproduce — or refute — your claim.

### 31.5.2 The Three Pillars of Reliable Benchmarking

Professional benchmarking rests on three foundational principles: repeatability, comparability, and completeness.

**Repeatability** means running the same experiment multiple times produces the same result. This requires fixed random seeds, consistent test datasets, and measuring variance across runs. A single measurement of “10.3ms” is worthless because you don’t know if that’s typical or an outlier. Measuring 100 times and reporting “10.0ms  $\pm$  0.5ms” tells you the true performance and its variability.

Here’s how your implementation ensures repeatability:

```
# Measure latency with statistical rigor
latencies = []
for _ in range(num_runs):
    start = time.time()
    _ = model.forward(X_test[:1]) # Single sample inference
    latencies.append((time.time() - start) * 1000) # Convert to ms

avg_latency = np.mean(latencies)
std_latency = np.std(latencies)
```

The loop runs inference 100 times (by default) to capture variance. The first few runs may be slower due to cold caches, and occasional runs may hit garbage collection pauses. By aggregating many measurements, you get a statistically valid estimate.

**Comparability** means different people can fairly compare results. This requires documenting the environment completely:

```
def _get_system_info(self):
    """Collect system information for reproducibility."""
    return {
        'platform': platform.platform(),
        'python_version': sys.version.split()[0],
        'numpy_version': np.__version__
    }
```

When someone sees your submission claiming 10ms latency, they need to know if that was measured on a MacBook or a server with 32 CPU cores. Platform differences can cause 10x performance variations, making cross-platform comparisons meaningless without context.

**Completeness** means capturing all relevant metrics, not cherry-picking favorable ones. Your `benchmark_model` method measures six distinct metrics:

```
self.metrics = {
    'parameter_count': int(param_count),
    'model_size_mb': float(model_size_mb),
    'accuracy': float(accuracy),
    'latency_ms_mean': float(avg_latency),
    'latency_ms_std': float(std_latency),
    'throughput_samples_per_sec': float(1000 / avg_latency)
}
```

Each metric answers a different question. Parameter count indicates model capacity. Model size determines deployment cost. Accuracy measures task performance. Latency affects user experience. Throughput determines batch processing capacity. Optimizations create trade-offs between these dimensions, so measuring all of them prevents hiding downsides.

### 31.5.3 End-to-End Inference Complexity

Before tuning any single knob, it helps to write down the whole inference cost in one place. For a transformer-style model serving a request at batch size  $B$ , sequence length  $S$ , hidden dimension  $D$ , and  $L$  layers, the end-to-end cost decomposes as:

$$T_{\text{inference}} = \underbrace{T_{\text{preproc}}}_{O(S)} + \underbrace{L \cdot (T_{\text{attn}} + T_{\text{ffn}})}_{\text{layer stack}} + \underbrace{T_{\text{postproc}}}_{O(1)}$$

- Attention per layer:  $O(B \cdot S^2 \cdot D)$  FLOPs,  $O(B \cdot S \cdot D)$  memory reads with KV caching,  $O(B \cdot S^2 \cdot D)$  without.
- Feed-forward per layer:  $O(B \cdot S \cdot D^2)$  FLOPs,  $O(D^2)$  weight bytes read per layer (amortized across  $B \cdot S$  tokens).
- Total:  $O(L \cdot B \cdot S \cdot D \cdot (S + D))$  FLOPs dominate training and prefill; at decode time,  $O(L \cdot B \cdot D^2)$  weight reads per generated token dominate — which is why decode is memory-bound.

Each optimization from Modules 14–18 attacks a specific term, and Module 19 verifies the combined effect:

- **Quantization** (15) shrinks the weight-byte constant  $4\times$  — wins on decode and embedding lookups.
- **Compression** (16) reduces the effective  $D$  (structured) or the nonzero count (unstructured) — wins on FFN cost.
- **Acceleration** (17) collapses memory complexity for element-wise ops ( $\Theta(kN) \rightarrow \Theta(N)$ ) and raises the GEMM constant toward peak.
- **Memoization** (18) turns the  $S^2$  attention term into  $S$  during decode — wins more as context grows.

A capstone submission must measure the *composite* effect. Optimizations interact — pruning that enables lower rank may hurt the quantization scale; kernel fusion that assumes dense tensors may defeat sparsity — and only an end-to-end benchmark reveals whether the stack actually compounds.



### 31.5.4 Latency vs Throughput: A Critical Distinction

Many beginners confuse latency and throughput because both relate to speed. They measure fundamentally different things.

**Latency** measures per-sample speed: how long does it take to process one input? This matters for real-time applications where users wait for results. Your implementation measures latency by timing single-sample inference:

```
# Latency: time for ONE sample
start = time.time()
_ = model.forward(X_test[:1]) # Shape: (1, features)
latency_ms = (time.time() - start) * 1000
```

A model with 10ms latency processes one input in 10 milliseconds. The user submits a query and waits 10ms. That number lives or dies on user experience.

**Throughput** measures batch capacity: how many inputs can you process per second? This is the metric for offline jobs grinding through millions of examples. Your implementation derives it from latency:

```
throughput_samples_per_sec = 1000 / avg_latency
```

At 10ms per sample, throughput is  $1000 / 10 = 100$  samples/second — but only if you process one sample at a time. Batching changes the math. A batch of 32 samples might take 50ms total, which is 640 samples/second of throughput at the cost of 50ms per-request latency.

The trade-off: **Batching increases throughput but hurts latency**. A production API serving individual user requests optimizes for latency to maintain a snappy user experience. A background batch processing pipeline optimizes for throughput to minimize total compute costs. Choosing between them is a fundamental architectural decision.

#### **i** Systems Implication: Hardware Utilization and Little's Law

To truly master this engineering trade-off, you must look beyond timing metrics and analyze Hardware Utilization. A GPU processing a single token (batch size 1) might only utilize 5% of its available ALUs, wasting 95% of the silicon you are paying for. As you increase the batch size, you saturate the computation engines and memory buses, pushing utilization toward 100%. However, **Little's Law** ( $L = \lambda W$ ) mathematically dictates that the number of concurrent items in the system ( $L$ ) equals the throughput ( $\lambda$ ) multiplied by the latency ( $W$ ). In a capstone setting—and in production—you must intentionally target a specific operating point on this curve. You cannot cheat the physical limits of the underlying hardware.

### 31.5.5 Statistical Rigor: Why Variance Matters

Single measurements lie. Variance tells the truth about performance consistency.

Consider two models, both with mean latency of 10.0ms. Model A has standard deviation of 0.5ms. Model B has standard deviation of 4.2ms. Which would you deploy?

Model A's predictable performance (9.5-10.5ms range) provides consistent user experience. Model B's erratic performance (sometimes 6ms, sometimes 15ms) creates frustration. Users prefer reliable slowness over unpredictable speed.

Your implementation captures this variance:

```
latencies = []
for _ in range(num_runs):
    start = time.time()
    _ = model.forward(X_test[:1])
    latencies.append((time.time() - start) * 1000)

avg_latency = np.mean(latencies)
std_latency = np.std(latencies) # Captures variance
```

Running 100 iterations isn't just for accuracy of the mean. It also characterizes the distribution. High standard deviation indicates performance varies significantly run-to-run, perhaps due to garbage collection pauses, cache effects, or OS scheduling.

In production systems, engineers track percentiles (p50, p90, p99) to understand tail latency. The p99 latency tells you “99% of requests complete within this time,” which matters more for user experience than mean latency. One user experiencing a 100ms delay (because they hit p99) has a worse experience than if all users consistently saw 20ms.

### 31.5.6 The Optimization Trade-off Triangle

Every optimization involves trade-offs between three competing objectives: speed (latency), size (memory), and accuracy. You can optimize for any two, but achieving all three simultaneously is impossible with current techniques.

**Fast + Small** means aggressive optimization. Quantizing to INT8 reduces model size 4x and speeds up inference 2x, but typically costs 1-2% accuracy. Pruning 50% of weights halves memory and adds another speedup, but may lose another 1-2% accuracy. You've traded accuracy for efficiency.

**Fast + Accurate** means careful optimization. You might quantize only certain layers, or use INT16 instead of INT8. You preserve accuracy but achieve less compression. The model is faster but not dramatically smaller.

**Small + Accurate** means conservative techniques. Knowledge distillation transfers accuracy from a large teacher to a small student. The student is smaller and maintains accuracy, but may be slower than aggressive quantization because it still operates in FP32.

Your submission captures these trade-offs automatically:

```
submission['improvements'] = {
    'speedup': float(baseline_latency / optimized_latency),
    'compression_ratio': float(baseline_size / optimized_size),
    'accuracy_delta': float(
        optimized_report.metrics['accuracy'] - baseline_report.metrics['accuracy']
    )
}
```

A speedup of 2.3x with compression of 4.1x but accuracy delta of -0.01 (-1%) shows you chose the “fast + small” corner of the triangle. A speedup of 1.2x with compression of 1.5x but accuracy delta of 0.00 shows you chose “accurate + moderately fast.”

### 31.5.7 Schema Validation: Enabling Automation

Your submission format uses a structured JSON schema that enforces completeness and type safety. This isn't bureaucracy—it enables powerful automation.

Without schema validation, submissions become inconsistent. One person reports accuracy as a percentage string (“92%”), another as a float (0.92), another as an integer (92). Aggregating these results requires manual cleaning. With schema validation, every submission uses the same format:

```
# Schema-enforced format
'accuracy': float(accuracy) # Always 0.0-1.0 float

# Validation catches errors
assert 0 <= metrics['accuracy'] <= 1, "Accuracy must be in [0, 1]"
```

This enables automated processing:

The code in `?@lst-20-capstone-schema-automation` makes this concrete.

```
# Aggregate community results automatically
all_submissions = [load_json(f) for f in submission_files]
avg_accuracy = np.mean([s['baseline']['metrics']['accuracy']
                        for s in all_submissions])

# Build leaderboards
sorted_by_speedup = sorted(all_submissions,
                           key=lambda s: s['improvements']['speedup'],
                           reverse=True)

# Detect regressions in CI/CD
if new_latency > baseline_latency * 1.1:
    raise Exception("Performance regression: 10% slower!")
```

: Listing 20.1 — Schema-validated submissions unlock automated aggregation, leaderboards, and regression detection. `{@lst-20-capstone-schema-automation}`

The schema also enables forward compatibility. When you add new optional fields, old submissions remain valid. When you require new fields, the version number increments, and validation enforces the migration.

### 31.5.8 Performance Measurement Traps

Real-world benchmarking is full of subtle traps that invalidate measurements. Understanding these pitfalls is crucial for accurate results.

**Trap 1: Measuring the Wrong Thing.** If you time model creation instead of just inference, you’re measuring initialization overhead, not runtime performance. If you include data loading in the timing loop, you’re measuring I/O speed, not model speed. The fix is isolating exactly what you want to measure:

The code in `?@lst-20-capstone-trap-isolation` makes this concrete.

```
# Prepare data BEFORE timing
X = create_test_input()

# Time ONLY the operation you care about
start = time.time()
output = model.forward(X) # Only this is timed
latency = time.time() - start
```

```
# Process output AFTER timing
predictions = postprocess(output)
```

: Listing 20.2 — Isolating the timed region from setup and post-processing. {#lst-20-capstone-trap-isolation}

**Trap 2: Ignoring System Noise.** Operating systems multitask. Your benchmark might get interrupted by background processes, garbage collection, or CPU thermal throttling. Single measurements capture noise. Multiple measurements average it out. Your implementation runs 100 iterations by default to handle this.

**Trap 3: Cold Start Effects.** The first inference is often slower because caches are cold and JIT compilers haven’t optimized yet. Production benchmarks typically discard the first N runs as “warm-up.” Your implementation includes warm-up inherently by averaging all runs—the few slow cold starts get averaged with many fast warm runs.

**Trap 4: Batch Size Confusion.** Measuring latency on `batch_size=32` then dividing by 32 doesn’t give per-sample latency. Batching amortizes overhead, so `batch latency / batch_size` underestimates per-sample latency. Always measure with the same batch size as production deployment.

### 31.5.9 System Integration: The Complete ML Lifecycle

This capstone represents the final stage of the ML systems lifecycle, but it’s also the beginning of the next iteration. Production ML systems operate in a never-ending loop:

1. **Research & Development** - Build models (Modules 01-13)
2. **Baseline Measurement** - Benchmark unoptimized performance (Module 19)
3. **Optimization** - Apply techniques like quantization and pruning (Modules 14-18)
4. **Validation** - Benchmark optimized version (Module 19)
5. **Decision** - Keep optimization if improvements outweigh costs (Module 20)
6. **Deployment** - Serve model in production
7. **Monitoring** - Track performance over time, detect regressions
8. **Iteration** - When performance degrades or requirements change, loop back to step 3

Your submission captures a snapshot of this cycle. The baseline metrics document performance before optimization. The optimized metrics show results after applying techniques. The improvements section quantifies the delta. The `techniques_applied` list enables reproducibility.

In production, engineers maintain this documentation across hundreds of experiments. When a deployment’s latency increases from 10ms to 30ms three months later, they consult the original benchmark to understand what changed. Without `system_info` and reproducible measurements, debugging becomes guesswork.

## 31.6 Production Context

### 31.6.1 Your Implementation vs. Industry Standards

Your TinyTorch benchmarking system implements the same principles used by production ML frameworks and research competitions, just at educational scale.

Table 31.4 places your implementation side by side with the production reference for direct comparison.

Table 31.4: Feature comparison between TinyTorch capstone benchmarking and industry systems.

| Feature           | Your Implementation                      | Production Systems                                  |
|-------------------|------------------------------------------|-----------------------------------------------------|
| <b>Metrics</b>    | 6 core metrics (accuracy, latency, etc.) | 20+ metrics including p99 latency, memory bandwidth |
| <b>Runs</b>       | 100 iterations for variance              | 1000+ runs, discard outliers                        |
| <b>Validation</b> | Python assertions                        | JSON Schema, automated CI checks                    |
| <b>Format</b>     | Simple JSON                              | Protobuf, versioned schemas                         |
| <b>Scale</b>      | Single model benchmarks                  | Automated pipelines tracking 1000s of experiments   |

### 31.6.2 Code Comparison

The following comparison shows how your educational implementation translates to production tools.

## 31.7 Your TinyTorch

```
from tinytorch.olympics import BenchmarkReport, generate_submission

# Benchmark baseline
baseline_report = BenchmarkReport(model_name="my_model")
baseline_report.benchmark_model(model, X_test, y_test, num_runs=100)

# Benchmark optimized
optimized_report = BenchmarkReport(model_name="optimized_model")
optimized_report.benchmark_model(opt_model, X_test, y_test, num_runs=100)

# Generate submission
submission = generate_submission(
    baseline_report=baseline_report,
    optimized_report=optimized_report,
    techniques_applied=["quantization", "pruning"]
)

save_submission(submission, "results.json")
```

## 31.8 Production MLflow

```
import mlflow

# Track baseline experiment
with mlflow.start_run(run_name="baseline"):
    mlflow.log_params({"model": "my_model"})
    metrics = benchmark_model(model, X_test, y_test)
    mlflow.log_metrics(metrics)
```

```

mlflow.log_artifact("model.pkl")

# Track optimized experiment
with mlflow.start_run(run_name="optimized"):
    mlflow.log_params({"model": "optimized_model",
                      "techniques": ["quantization", "pruning"]})
    metrics = benchmark_model(opt_model, X_test, y_test)
    mlflow.log_metrics(metrics)
    mlflow.log_artifact("optimized_model.pkl")

# Compare experiments in MLflow UI

```

Let's walk through the comparison line by line:

- **Line 1 (Import):** TinyTorch uses a simple module import. MLflow provides enterprise-grade experiment tracking with databases and web UIs.
- **Line 4 (Benchmark baseline):** TinyTorch's `BenchmarkReport` mirrors MLflow's experiment runs. Both capture metrics and system context.
- **Line 8 (Benchmark optimized):** Same API in both—create report, benchmark model. This consistency makes transitioning to production tools natural.
- **Line 12 (Generate submission):** TinyTorch generates JSON. MLflow logs to a database that supports querying, visualization, and comparison.
- **Line 18 (Save):** TinyTorch saves to file. MLflow persists to SQL database with version control and artifact storage.

#### 💡 What's Identical

The workflow pattern: baseline → optimize → benchmark → compare → decide. Whether you use TinyTorch or MLflow, this cycle is fundamental to production ML. The tools scale, but the methodology remains the same.

### 31.8.1 Why Benchmarking Matters at Scale

Benchmarking earns its keep when the numbers get large:

- **Model serving.** A recommendation system handles 10M requests/day. Cutting latency from 20ms to 10ms saves 100,000 seconds of compute every day — 1.16 days of CPU time daily, a 50% serving-cost reduction. That's a budget line, not a metric.
- **Training efficiency.** A \$1M LLM training run is 60% bottlenecked on data loading. Fixing the input pipeline frees \$600,000 of GPU time. You only know to look because the profiler said so.
- **Deployment constraints.** A mobile app must fit a 50MB model budget. Quantization shrinks a 200MB checkpoint to 50MB at the cost of 1% accuracy. The app ships — and the only reason anyone signed off on the trade is because the trade was measured.

Reproducible benchmarking isn't academic. It's how engineers defend a decision in a design review and prove an optimization paid off in a postmortem.

## 31.9 Check Your Understanding

### 💡 Check Your Understanding — Capstone

Before moving on, verify you can articulate each of the following:

- Why an MLPerf-style submission requires reproducibility infrastructure — system metadata, schema validation, versioned artifacts — not just a trained model.
- How end-to-end inference complexity decomposes across attention, FFN, and decode terms, which optimization from Modules 14–18 attacks each, and how Module 19 validates the combined effect.
- The three pillars — repeatability, comparability, completeness — and what fails when any one is missing.
- Why compounded optimizations (prune → quantize → fuse → cache) must be measured end-to-end, since their interactions can defeat each other.

If any of these feels fuzzy, revisit the Core Concepts section (especially The Three Pillars of Reliable Benchmarking and Schema Validation) before moving on.

Test yourself with these systems thinking questions about benchmarking and performance measurement.

#### Q1: Memory Calculation

A model has 5 million parameters stored as FP32. After INT8 quantization, how much memory is saved?

### 💡 Answer

FP32:  $5,000,000 \text{ params} \times 4 \text{ bytes} = 20,000,000 \text{ bytes} = \mathbf{19.07 \text{ MB}}$

INT8:  $5,000,000 \text{ params} \times 1 \text{ byte} = 5,000,000 \text{ bytes} = \mathbf{4.77 \text{ MB}}$

Savings:  $19.07 \text{ MB} - 4.77 \text{ MB} = \mathbf{14.30 \text{ MB}}$  (75% reduction).

Compression ratio:  $19.07 / 4.77 = \mathbf{4.0\times}$ .

This is why quantization is standard for mobile deployment — the model has to fit the budget.

#### Q2: Latency Variance Analysis

Model A:  $10.0\text{ms} \pm 0.3\text{ms}$  latency. Model B:  $10.0\text{ms} \pm 3.0\text{ms}$  latency. Same accuracy. Which do you deploy and why?

### 💡 Answer

#### Deploy Model A.

Same mean (10.0ms), but Model A has 10× lower variance (0.3ms vs 3.0ms std).

- Model A's 95% range ( $\pm 2\sigma$ ): ~9.4–10.6ms
- Model B's 95% range ( $\pm 2\sigma$ ): ~4.0–16.0ms

#### Why consistency wins:

- Users tolerate predictable slowness; they don't tolerate jitter.
- High variance is usually GC pauses, cache misses, or resource contention — three problems you don't want in your hot path.
- Production SLAs are written against p99, not mean. Model B's p99 could land near 16ms while Model A's stays under 11ms.

In production, **reliability beats average speed**. A consistently decent experience trumps an unreliably fast one.

#### Q3: Batch Size Trade-off

Measuring latency with `batch_size=32` gives 100ms total. Can you claim  $100 / 32 = 3.1\text{ms}$  per-sample latency?

#### 💡 Answer

**No.** Amortized latency is not real latency.

Batching amortizes fixed overhead (data transfer, kernel launch). Per-sample latency at `batch=1` is higher than the batch-32 number divided by 32.

What's actually going on:

- Batch=32: 100ms total  $\rightarrow$  3.1ms/sample (amortized)
- Batch=1: 8ms total  $\rightarrow$  8ms/sample (actual)

The cost decomposes as:

- Fixed overhead:  $\sim 5.0\text{ms}$  (data transfer, kernel launch)
- Variable cost:  $\sim 95.0\text{ms} / 32 = \sim 3.0\text{ms}$  per sample
- At `batch=1`:  $5.0\text{ms} + 3.0\text{ms} \approx 8.0\text{ms}$

**Always benchmark at the deployment batch size.** If production serves single requests, measure with `batch=1`.

#### Q4: Speedup Calculation

Baseline: 20ms latency. Optimized: 5ms latency. What is the speedup and what does it mean?

#### 💡 Answer

Speedup = baseline / optimized =  $20\text{ms} / 5\text{ms} = 4.0\times$ .

What that buys you:

- Same input processed in one-quarter the time.
- Same hardware now serves  $4\times$  the traffic.

What it means in dollars and SLA:

- 100 req/sec baseline  $\rightarrow$  400 req/sec optimized.
- \$1,000/month compute  $\rightarrow$  \$250/month for the same load.
- 60% utilization at SLA  $\rightarrow$  85% headroom before you have to scale.

**Caveat:** speedup is one number. Always read it alongside `accuracy_delta` and `compression_ratio` — a  $4\times$  speedup that costs 5% accuracy is not the same product.

#### Q5: Schema Validation Value

Why does the submission schema require `accuracy` as float in `[0, 1]` instead of allowing any format?

#### 💡 Answer

**Type safety enables automation.**

Without schema:

```
# Different submissions, different formats (breaks aggregation)
{"accuracy": "92%"}           # String
{"accuracy": 92}              # Integer (out of 100)
{"accuracy": 0.92}           # Float
{"accuracy": "good"}          # Non-numeric
```



Aggregating these requires manual parsing and error handling.  
With schema:

```
# All submissions use same format (aggregation works)
{"accuracy": 0.92} # Always float in [0.0, 1.0]
```

**Benefits:** 1. **Automated validation** - Reject invalid submissions immediately 2. **Aggregation** - `np.mean([s['accuracy'] for s in submissions])` just works 3. **Comparison** - Sort by accuracy without parsing different formats 4. **APIs** - Other tools can consume submissions without custom parsers

**Real example:** Papers with Code leaderboards require strict schemas. Thousands of submissions from different teams aggregate automatically because everyone follows the same format.

## 31.10 Key Takeaways

- **A submission is an artifact, not a number:** reproducibility infrastructure — system metadata, schema-validated JSON, versioned code — is what makes a claim defensible a year later.
- **End-to-end measurement beats per-optimization microbenchmarks:** optimizations interact, and only the composite benchmark reveals whether the stack compounds or collides.
- **Three pillars are non-negotiable:** repeatability (many runs, variance reported), comparability (platform recorded), completeness (all six metrics, no cherry-picks).
- **The trade-off triangle is real:** fast + small + accurate is not simultaneously achievable with current techniques; every submission picks a corner, and the schema forces you to own the choice.

**Coming next:** the framework is finished. The Conclusion chapter steps back to review which abstractions paid off, which trade-offs recur, and how what you built compares to the production systems that inspired it — read it with your `results.json` open.

## 31.11 Further Reading

For students who want to understand the academic foundations and industry standards for ML benchmarking:

### 31.11.1 Seminal Papers

- **MLPerf: An Industry Standard Benchmark Suite for Machine Learning Performance** - Mattson et al. (2020). Defines standardized ML benchmarks for hardware comparison. The gold standard for fair performance comparisons. **Systems Implication:** Forced hardware vendors to transparently report power consumption and system topology alongside raw FLOPS, cementing the reality that interconnects and power walls dictate cluster-scale ML. [arXiv:1910.01500](#)
- **A Step Toward Quantifying Independently Reproducible Machine Learning Research** - Pineau et al. (2021). Analyzes reproducibility crisis in ML and proposes requirements for verifiable claims. Introduces reproducibility checklist adopted by NeurIPS. **Systems Implication:** Emphasized that identical code can yield divergent results across different hardware architectures (e.g., due to floating-point non-determinism in parallel reductions), making hardware logging a strict requirement for reproducibility. [arXiv:2104.05563](#)
- **Hidden Technical Debt in Machine Learning Systems** - Sculley et al. (2015). Identifies systems challenges in production ML, including monitoring, versioning, and reproducibility. Required reading for ML systems engineers. **Systems Implication:** Revealed that the underlying compute infrastructure,

data pipelines, and serving systems dwarf the actual ML code, showing that scaling ML is fundamentally a distributed systems engineering problem. [NeurIPS 2015](#)

### 31.11.2 Additional Resources

- **MLflow Documentation:** <https://mlflow.org/> - Production experiment tracking system implementing patterns from this module
- **Papers with Code:** <https://paperswithcode.com/> - See how research papers submit benchmarks with reproducible code
- **Weights & Biases Best Practices:** <https://wandb.ai/site/experiment-tracking> - Industry standard for ML experiment management

## 31.12 What's Next

**i** The framework is finished. The learning isn't.

Twenty modules ago, `Tensor` was an empty class. Now it has autograd behind it, a transformer on top of it, a quantizer that compresses it, and a benchmark report that proves what it can do. You didn't read about how PyTorch is built. You built it.

The `results.json` you ship from this module is the first artifact in a long career of them. Every system you'll deploy from here forward gets the same treatment: baseline, optimize, measure, justify.

Where to take the framework next:

Table 31.5 collects suggested directions for extending the framework beyond the capstone.

Table 31.5: **Suggested directions for extending the framework beyond the capstone.**

| Direction                            | What to build                                                                             | What it teaches                                                    |
|--------------------------------------|-------------------------------------------------------------------------------------------|--------------------------------------------------------------------|
| <b>Push the optimizers</b>           | Benchmark milestone models (TinyDigits CNN, Transformer) end-to-end through Modules 14-18 | How real optimizations interact under measurement                  |
| <b>Scale to a team</b>               | Wire MLflow or Weights & Biases into the same <code>BenchmarkReport</code> flow           | How professional experiment tracking grows from this same skeleton |
| <b>Publish results</b>               | Convert your schema into a Papers with Code submission                                    | How reproducibility becomes a community contract                   |
| <b>Automate regression detection</b> | Run the benchmark on every commit in CI                                                   | How performance is defended, not just achieved                     |

The next chapter — the **Conclusion** — steps back from the code. It looks at what you actually learned by writing all of it: which abstractions paid off, which trade-offs you'll keep meeting, and how this framework you built compares to the production systems that inspired it. Read it with your `results.json` open.

## 31.13 Get Started

### Interactive Options

- **Launch Binder** - Run interactively in browser, no setup required
- **View Source** - Browse the implementation code

### Save Your Progress

Binder sessions are temporary. Download your completed notebook when done, or clone the repository for persistent local work.



## Chapter 32

# You Built Something Real

---

At the start of this journey, we made a simple promise: **Don't import torch. Build it.** You did.

## 32.1 What You Accomplished

Across 20 modules and six historical milestones, you built a complete machine learning framework from scratch:

### Foundation Tier (Modules 01-08)

- Tensors with broadcasting and shape manipulation
- Activation functions with gradients
- Linear layers with proper initialization
- Loss functions with numerical stability
- DataLoaders with batching and shuffling
- Automatic differentiation with computation graphs
- Optimizers (SGD, Adam, RMSprop)
- Complete training loops

### Architecture Tier (Modules 09-13)

- Convolutional layers (Conv2d, MaxPool2d)
- Tokenization for text processing
- Embeddings (token and positional)
- Multi-head self-attention
- Transformer blocks with LayerNorm

### Optimization Tier (Modules 14-19)

- Profiling and bottleneck identification
- Quantization (INT8, FP16)
- Model compression and pruning
- Acceleration techniques
- KV-cache for generation speedup
- Benchmarking infrastructure

Then you proved it works. You recreated six decades of neural network breakthroughs, from the 1958 Perceptron to 2018 MLPerf optimization, all running on code you wrote.

## 32.2 The Mindset Shift

Something changed along the way. When you started, `import torch` was magic. Now you know:

- **Tensors** are not mysterious. They are multidimensional arrays with broadcasting rules you implemented.
- **Autograd** is not a black box. It is a computation graph you built and traversed.

- **Attention** is not incomprehensible. It is matrix multiplication with learned weights you coded.
- **Optimization** is not guesswork. It is systematic measurement and targeted improvement you executed.

You went from user to builder. From “it works somehow” to “I know exactly how it works.”

### 32.3 What You Can Do Now

That foundation gives you four capabilities you didn’t have before.

#### Debug Production Issues

When a model runs out of memory, you understand tensor allocation. When gradients explode, you can trace the computation graph. When training is slow, you know where to profile. You built these systems. You can fix them.

#### Read Framework Source Code

PyTorch’s `torch.nn.Linear` follows the same architecture as your Module 03 implementation. The autograd engine uses the same topological sort you wrote. The patterns are familiar because you built them first at educational scale.

#### Optimize for Deployment

You know that quantization trades precision for memory. You know that pruning removes parameters without destroying accuracy. You know that KV-caching speeds up generation. These are not abstract concepts. They are techniques you implemented and measured.

#### Contribute to Open Source

The gap between TinyTorch and production frameworks is scale and optimization, not architecture. You understand the design. Contributing a new layer, optimizer, or feature is extending patterns you already know.

### 32.4 Your Code vs Production Frameworks

Your TinyTorch implementation and PyTorch share the same core architecture:

Table 32.1 places your implementation side by side with the production reference for direct comparison.

Table 32.1: Component-by-component comparison of TinyTorch and PyTorch internals.

| Component  | Your TinyTorch                   | PyTorch                          | The Difference                     |
|------------|----------------------------------|----------------------------------|------------------------------------|
| Tensor     | Pure Python, NumPy backend       | C++/CUDA, optimized memory       | Performance, not architecture      |
| Autograd   | Python computation graph         | C++ tape-based                   | Same algorithm, different language |
| Layers     | Module pattern, forward/backward | Module pattern, forward/backward | Nearly identical API               |
| Optimizers | State dict, step method          | State dict, step method          | Same interface                     |
| Attention  | QKV projection, softmax, output  | QKV projection, softmax, output  | Same math                          |

The principles transfer directly. What you learned scales.

## 32.5 Paths Forward

Your TinyTorch foundation opens four directions:

### Research

Implement new architectures from papers. You understand the building blocks. Novel attention mechanisms, new normalization techniques, experimental optimizers: these are combinations of components you already built.

### Production ML Engineering

Apply optimization techniques to real systems. Profile before optimizing. Quantize for deployment. Cache for inference speed. Production teams need these skills.

### Framework Development

Contribute to PyTorch, TensorFlow, JAX, or emerging frameworks. You understand their architecture because you built a working version. Contributing is easier when you already know how the pieces fit together.

### Teaching

Use TinyTorch to teach others. The progression from tensors to transformers is a curriculum you already taught yourself. Help the next generation of builders see what lives inside the black box.

## 32.6 The Broader Mission

TinyTorch is part of the [Machine Learning Systems](#) project, an open effort to train the next generation of ML systems engineers. You are now part of a community of builders who chose to understand deeply rather than use superficially.

The world has enough users. It needs more builders: people who can debug, optimize, adapt, and extend systems when the abstractions break down. You chose to become one.

## 32.7 A Final Note

In the preface, we wrote:

Everyone wants to be an astronaut. Very few want to be the rocket scientist.

You chose to be both. From your first perceptron to your final transformer, you wrote every line. You traced every gradient. You watched the loss curve bend because of code you shipped. You did not just fly the rocket — you built it. And now you understand why it flies.

**Don't import torch. You built it.**

### 💡 A note from your instructor

If you made it this far, I want to say something directly.

This lab guide was written to be the thing I wish had existed when I was learning how these systems actually work. Not a tutorial. Not a survey. A *build book* — where every abstraction eventually ends at code you wrote yourself.

You chose the harder path. Twenty modules, six milestones, several hundred pages, and more than a few gnarly bugs between you and here. Most people don't finish. Most people read the textbook, watch the videos, import the library, and call it understanding. You went further.

The framework you built is small. The mental model you built is not. Carry it with you. The next time you read a paper, use a new architecture, or debug a system that's slower than it should be — you'll know where to look. You'll know what to ask. That is what this guide was for.

Thank you for building TinyTorch with me.

— Vijay Janapa Reddi  
Harvard University



 Chapter 33

# Glossary

---

Alphabetical reference for the key technical terms introduced across the TinyTorch Lab Guide, framed in the voice the modules use them in. Where a term is central to a specific module, that module is cited as the place to return for depth.

## 33.1 A

**Acceleration.** The practice of replacing slow Python loops with vectorized NumPy, cache-aware memory layouts, and BLAS-backed kernels so a computation runs closer to the hardware’s peak throughput (Module 17).

**Activation Function.** An element-wise nonlinearity such as ReLU, Sigmoid, Tanh, or GELU applied between linear layers; without it, a stack of layers collapses algebraically into a single linear transform (Module 02).

**Activation Pinning.** The systems consequence of reverse-mode autograd: every forward activation must stay resident in memory until the backward pass consumes it, which sets a hard floor on training VRAM (Module 06).

**Adam.** An adaptive optimizer that maintains per-parameter first and second moment estimates of the gradient, trading extra optimizer state for faster, more robust convergence (Module 07).

**AdamW.** Adam with decoupled weight decay, where the L2 penalty is applied directly to the parameter rather than folded into the gradient; the default optimizer for modern transformer training (Module 07).

**AllReduce.** A collective communication primitive that sums tensors across devices and returns the result to all participants; inserted by gradient clipping and synchronous data-parallel training, and often the step that bottlenecks distributed pipelines.

**Arithmetic Intensity.** The ratio of floating-point operations performed per byte of memory moved; the coordinate along the roofline model’s x-axis that determines whether a kernel is memory-bound or compute-bound (Module 17).

**Attention.** A mechanism that computes a weighted combination of value vectors using similarities between a query and a set of keys, enabling a model to look at every position in a sequence in parallel (Module 12).

**Autograd.** The automatic differentiation engine that records a computational graph during the forward pass and traverses it in reverse to produce gradients with respect to every parameter (Module 06).

**Autoregressive Generation.** Producing one token at a time, feeding each output back in as the next input; the decoding pattern that makes the KV cache a necessity rather than a nicety (Module 13).

## 33.2 B

**Backward Pass.** The traversal of the computational graph from loss to inputs that applies the chain rule to accumulate gradients into every parameter tensor (Module 06).

**Batch.** A group of examples processed together so a single matrix multiply can amortize memory-bandwidth and kernel-launch overhead across many samples (Module 05).

**Batch Normalization.** A layer that normalizes activations using batch statistics during training and running statistics during inference, stabilizing deep networks at the cost of a train/eval mode distinction (Module 09).

**Benchmarking.** The disciplined measurement of latency, throughput, and memory under controlled conditions, with warmup, multiple runs, and variance reporting, so that a reported speedup is reproducible (Module 19).

**BF16.** Brain Floating Point 16, a 16-bit format that keeps FP32's exponent range but halves the mantissa, trading precision for memory bandwidth during training (Module 15).

**BLAS.** Basic Linear Algebra Subprograms, the decades-old library interface that underlies every production matmul; its Level 3 routines are tiled and cache-aware so dense linear algebra approaches peak FLOPS (Module 17).

**Broadcasting.** The right-to-left shape-alignment rule that lets a  $(768, )$  bias add into a  $(32, 512, 768)$  activation tensor without materializing the tiled copy, saving memory proportional to the broadcast dimensions (Module 01).

**Byte Pair Encoding (BPE).** A subword tokenization algorithm that greedily merges the most frequent adjacent pairs, compressing text into a fixed vocabulary that stays robust to rare words (Module 10).

### 33.3 C

**Cache Tiling.** Breaking a matmul into blocks sized to fit in L1 or L2 cache so each loaded tile is reused many times before eviction, turning an  $O(n^3)$  operation from memory-bound back into compute-bound (Module 01).

**Calibration.** Passing a small representative dataset through a model to collect the activation statistics that post-training quantization needs to pick scale and zero-point parameters (Module 15).

**Causal Mask.** An upper-triangular mask applied to attention scores that zeroes out future positions, enforcing the autoregressive constraint that token  $t$  cannot attend to token  $t + 1$  (Module 12).

**Chain Rule.** The calculus identity that lets autograd compose per-operation local gradients into a full end-to-end gradient, one multiplication per edge of the computational graph (Module 06).

**Checkpointing.** Writing optimizer state, model weights, and training metadata to disk at regular intervals so a multi-day run can survive a preemption or crash (Module 08).

**Compression Ratio.** The size of the original model divided by the size after pruning, distillation, or low-rank approximation; the headline number every deployment constraint is measured against (Module 16).

**Computational Graph.** The directed acyclic graph of tensor operations autograd records on the forward pass; its topology determines the order and memory cost of the backward pass (Module 06).

**Compute-Bound.** A kernel whose runtime is limited by the rate at which the ALUs can issue floating-point operations, not by the rate at which memory can feed them (Module 14).

**Contiguous Layout.** A tensor whose elements sit in memory in the order implied by its shape and row-major strides; iteration is sequential, cache lines are reused, and SIMD loads stay aligned (Module 01).

**Convergence.** The state in which additional optimizer steps no longer meaningfully reduce training loss; the signal that a run is either done or stuck (Module 07).

**Convolution.** A weight-sharing layer that slides a small kernel across spatial input, exploiting the locality and translation-equivariance of image data (Module 09).

**Cross-Entropy Loss.** The negative log-likelihood of the true class under the model's predicted distribution, the standard training signal for classification and language modeling (Module 04).

### 33.4 D

**DataLoader.** An iterator that streams minibatches from a dataset, optionally shuffling, augmenting, and prefetching so the accelerator is never idle waiting for input (Module 05).

**Dropout.** A regularizer that zeroes a random subset of activations during training and rescales the remainder, preventing co-adaptation; active only in train mode (Module 03).

**Dtype.** The numerical type of a tensor's elements (float32, float16, bfloat16, int8); the single knob that most directly trades precision for memory and bandwidth (Module 01).

### 33.5 E

**Embedding Table.** A learnable matrix of shape `(vocab_size, d_model)` whose rows map discrete token indices to dense vectors, the input side of every modern language model (Module 11).

**Epoch.** One full pass of the optimizer over every example in the training set; the unit most schedules (learning rate, weight decay) are expressed in (Module 08).

### 33.6 F

**Forward Pass.** The left-to-right evaluation of a model that turns an input batch into predictions and, during training, records the computational graph for the backward pass (Module 01).

**FP16.** Half-precision floating point, a 16-bit format with reduced exponent range; halves memory and doubles bandwidth relative to FP32, but requires loss scaling to avoid gradient underflow (Module 15).

### 33.7 G

**GEMM.** General Matrix-Matrix Multiplication, the BLAS Level 3 routine that every linear layer's forward and backward passes ultimately dispatch to (Module 03).

**Gradient.** The partial derivative of the loss with respect to a parameter; the signal the optimizer consumes to decide which way and how far to step (Module 06).

**Gradient Accumulation.** Summing gradients from several micro-batches into a parameter's `.grad` buffer before calling `optimizer.step()`, simulating a larger batch on memory-constrained hardware (Module 06).

**Gradient Checkpointing.** Discarding selected forward activations and recomputing them during the backward pass, trading compute for peak memory during training (Module 18).

**Gradient Clipping.** Rescaling the gradient vector so its global norm stays below a threshold, preventing occasional outlier batches from blowing up the parameter update (Module 08).

### 33.8 H

**HBM.** High Bandwidth Memory, the stacked DRAM directly attached to a GPU package; far faster than host DRAM but far slower than on-chip SRAM, which is why FlashAttention exists.

### 33.9 I

**INT8.** An 8-bit integer representation used for quantized weights and activations; reduces model size 4x versus FP32 and lets inference hardware use integer ALUs that are cheaper and faster than floating-point units (Module 15).

**In-place Operation.** An operator that overwrites its input buffer rather than allocating a new one, saving memory bandwidth at the cost of breaking autograd whenever the input is still needed for the backward pass (Module 02).

### 33.10 K

**Kernel Fusion.** Combining several element-wise operations into a single pass over memory so the activation is read and written once instead of once per op; one of the largest wins available to a framework (Module 17).

**Knowledge Distillation.** Training a smaller student network to match the output distribution of a larger teacher, compressing the teacher’s learned behavior into a model with fewer parameters (Module 16).

**KV Cache.** The memoization of key and value projections across autoregressive steps, converting attention’s per-step cost from  $O(N^2)$  recomputation to  $O(N)$  reuse at the cost of  $O(N)$  additional memory per token (Module 18).

### 33.11 L

**Latency.** The time to complete a single request end to end; the metric that governs user-facing interactive inference (Module 19).

**Layer Norm.** Per-example normalization across the feature dimension, the stabilizing ingredient that lets transformers stack to hundreds of layers without exploding or collapsing activations (Module 13).

**Learning Rate.** The scalar multiplier on the gradient in each optimizer step; the single hyperparameter most likely to decide whether a run converges at all (Module 07).

**Linear Layer.** A parameterized affine transformation  $y = xW^\top + b$ , the building block every deep network is ultimately assembled from (Module 03).

**Little’s Law.** The queueing identity  $L = \lambda W$  relating concurrency, throughput, and latency; the tool for reasoning about when a serving system is bandwidth-bound versus latency-bound (Modules 19, 20).

**Logits.** The raw unnormalized scores a classifier emits before Softmax; the input most numerically stable loss functions expect (Module 04).

**Loss.** A scalar measure of how wrong a model’s predictions are on a batch; the quantity the optimizer minimizes (Module 04).

**Low-Rank Approximation.** Factoring a weight matrix  $W$  as  $UV^\top$  with intermediate rank  $r \ll \min(m, n)$ , cutting parameter count and compute while preserving most of the matrix’s dominant behavior (Module 16).

### 33.12 M

**Matmul.** Matrix multiplication; the  $O(n^3)$  compute,  $O(n^2)$  memory operation that dominates training and inference time and that every systems optimization ultimately defends (Module 01).

**Memoization.** Caching the result of a deterministic computation so later calls with the same inputs become a table lookup instead of a recompute; the basis for the KV cache (Module 18).

**Memory-Bound.** A kernel whose runtime is limited by the rate at which memory can deliver operands, not by the rate at which the ALUs can consume them; the regime where arithmetic intensity is low (Module 14).

**Memory Wall.** The widening gap between compute throughput and memory bandwidth that makes data movement, not math, the dominant cost of modern ML workloads (Module 15).

**Mini-batch.** A subset of the training set processed in a single forward/backward step; a compromise between the variance of stochastic updates and the memory cost of full-batch gradient descent (Module 05).

**MLP.** A multi-layer perceptron: alternating linear layers and nonlinearities; the feedforward sublayer inside every transformer block (Module 13).

**Multi-Head Attention.** Running several attention heads in parallel over independently projected query, key, and value subspaces, then concatenating; lets the model attend to multiple relational patterns at once (Module 12).

### 33.13 O

**Optimizer.** The stateful object that consumes gradients and produces parameter updates, from plain SGD to Adam to AdamW (Module 07).

**Overfitting.** When a model’s training loss keeps decreasing while its validation loss rises; the signal that the network has started memorizing instead of generalizing (Module 08).

### 33.14 P

**PagedAttention.** An inference-time memory manager that stores the KV cache in fixed-size blocks and indirections through a page table, eliminating the fragmentation that naive contiguous caches suffer under variable sequence lengths (Module 18).

**Positional Encoding.** A signal added to token embeddings so attention, which is otherwise permutation-invariant, can distinguish position; sinusoidal, learned, and rotary variants all serve the same purpose (Module 11).

**Post-Training Quantization.** Converting a trained FP32 model to a lower-precision format (typically INT8) using a small calibration set, without retraining (Module 15).

**Profiling.** Measuring where a program actually spends its time and memory, by instrumenting regions with timers and allocators, so optimization effort is directed at the actual bottleneck (Module 14).

**Pruning.** Zeroing out weights deemed unimportant (by magnitude, by gradient, by structured group) to reduce the effective parameter count of a trained model (Module 16).

### 33.15 Q

**Q/K/V.** The query, key, and value projections attention produces from the input; Q asks “what am I looking for,” K advertises “what do I offer,” V carries “what I will hand back” (Module 12).

**Quantization.** Replacing FP32 tensors with a lower-precision integer or narrow-float representation, reducing memory, bandwidth, and often compute cost at a small accuracy price (Module 15).

### 33.16 R

**ReLU.** The rectified linear unit  $\max(0, x)$ ; cheap, non-saturating, and the activation that made deep networks trainable at scale (Module 02).

**Residual Connection.** A shortcut that adds a layer’s input to its output, giving the gradient a direct path to earlier parameters and making deep stacks optimizable (Module 13).

**Retain Grad.** The opt-in flag that tells autograd to keep a non-leaf tensor’s gradient around after the backward pass, at the cost of the memory autograd normally reclaims (Module 06).

**Roofline Model.** A performance chart plotting achievable FLOP/s against arithmetic intensity, with a memory-bandwidth ceiling on the left and a peak-compute ceiling on the right; tells you which wall your kernel is hitting (Module 14).

### 33.17 S

**Scale and Zero-Point.** The two parameters that map a floating-point range  $[x_{\min}, x_{\max}]$  into an integer range during quantization: scale sets the step size, zero-point sets the integer that represents zero (Module 15).

**SGD.** Stochastic Gradient Descent, the baseline optimizer: subtract the learning rate times the gradient from each parameter (Module 07).

**SIMD.** Single Instruction, Multiple Data; CPU vector instructions that apply one op to a register full of values in a single clock, the reason element-wise NumPy is orders of magnitude faster than a Python loop (Module 01).

**Softmax.** The exponentiate-and-normalize function that turns a vector of logits into a probability distribution; numerically stabilized in practice by subtracting the max before exponentiating (Modules 02, 12).

**Stride.** The number of bytes (or elements) to skip in memory to move one step along a given dimension; the per-axis metadata that makes transpose an  $O(1)$  view and lets reshape stay zero-copy when contiguous (Module 01).

### 33.18 T

**Tensor.** The N-dimensional array abstraction every module in the book operates on: data buffer plus shape, dtype, and stride metadata (Module 01).

**Throughput.** The number of requests or tokens a system processes per unit time; the metric that governs batch-oriented training and offline inference (Module 19).

**Tokenization.** Mapping raw text to a sequence of integer token IDs a model can consume; the interface between the language and the tensors (Module 10).

**Train/Eval Mode.** The toggle on the model that switches dropout, batch norm, and other stochastic or batch-statistical layers between their training and inference behaviors (Module 08).

**Training Loop.** The outer program that iterates over epochs and batches, running forward, loss, backward, and optimizer step in sequence and periodically logging, checkpointing, and evaluating (Module 08).

**Transformer Block.** The residual sandwich of layer norm, multi-head attention, layer norm, and MLP that stacks N times to form a GPT or BERT-style model (Module 13).

### 33.19 V

**Vectorization.** Expressing a computation as a whole-array operation so NumPy can dispatch it to SIMD or BLAS instead of interpreting a Python loop element by element (Module 17).

**View.** A tensor that shares its data buffer with another tensor and only differs in shape or stride metadata; zero-copy,  $O(1)$  to construct, and silently aliased with its source (Module 01).

**Vocabulary.** The finite set of tokens a tokenizer knows about, with one integer ID per entry; its size sets the width of the input embedding table and the output logit layer (Module 10).