




Introduction to

Machine Learning Systems

Vijay
Janapa Reddi

Introduction to Machine Learning Systems

From Algorithms to Silicon: The Engineering of Deep Learning Systems

Introduction to Machine Learning Systems

From Algorithms to Silicon: The Engineering of Deep Learning Systems

Vijay Janapa Reddi

The MIT Press
Cambridge, Massachusetts
London, England

The MIT Press
Massachusetts Institute of Technology
77 Massachusetts Avenue
Cambridge, MA 02139
mitpress.mit.edu

© 2027 Massachusetts Institute of Technology

This work is subject to a Creative Commons CC-BY-NC-ND license.

This license applies only to the work in full and not to any components included with permission. Subject to such license, all rights are reserved. No part of this book may be used to train artificial intelligence systems without permission in writing from the MIT Press.



The MIT Press would like to thank the anonymous peer reviewers who provided comments on drafts of this book. The generous work of academic experts is essential for establishing the authority and quality of our publications. We acknowledge with gratitude the contributions of these otherwise uncredited readers.

This book was prepared and formatted in Quarto by Željko Hrček. Printed and bound in the United States of America.

Library of Congress Cataloging-in-Publication Data is available.

ISBN: 978-0-262-05888-9

EU Authorised Representative: Easy Access System Europe, Mustamäe tee 50, 10621 Tallinn, Estonia | Email: gpsr.requests@easproject.com

*In loving memory of my father,
whose lifelong passion for learning
and unwavering commitment to quality
still challenge me daily to strive for excellence.*

*And with deepest gratitude to my mother,
whose everyday generosity taught me,
more by example than instruction,
that anything worth having
is worth sharing.*

Contents

Foreword	xiii
Author’s Note	xvii
Preface	xix
Acknowledgments	xxv
A Note on AI Assistance	xxvii
Notation	xxix

Main Text

Chapter 1 Introduction	3
1.1 The AI Systems Moment	4
1.2 Data-Centric Paradigm Shift	4
1.3 The Evolution of AI Bottlenecks	6
1.4 The Bitter Lesson	10
1.5 Defining ML Systems	11
1.6 ML vs. Traditional Software	16
1.7 Iron Law of ML Systems	17
1.8 Three Dimensions of ML Efficiency	21
1.9 AI Engineering as a Discipline	24
1.10 Five-Pillar Framework	28
1.11 Book Organization	30
1.12 Fallacies and Pitfalls	31
1.13 Summary	32

Part I: Foundations of ML Systems

Foundation Principles	37
Chapter 2 ML Systems	39
2.1 Deployment Paradigm Framework	40
2.2 Physical Constraints: Why Paradigms Exist	41
2.3 Analyzing Workloads	43
2.4 System Balance and Hardware	47
2.5 Cloud ML: Computational Power	50
2.6 Edge ML: Latency and Privacy	55
2.7 Mobile ML: Offline Intelligence	61
2.8 TinyML: Ubiquitous Sensing	64
2.9 Paradigm Selection	67
2.10 Hybrid Architectures	71

2.11	System Entropy: Why Deployment Is Not the End	75
2.12	Fallacies and Pitfalls	75
2.13	Summary	77
Chapter 3	ML Workflow	79
3.1	ML Lifecycle	80
3.2	Lifecycle Stages	84
3.3	Problem Definition	88
3.4	Data Collection	90
3.5	Model Development	93
3.6	Evaluation and Validation	96
3.7	Deployment and Integration	98
3.8	Monitoring and Maintenance	100
3.9	Systems Thinking	102
3.10	Fallacies and Pitfalls	104
3.11	Summary	105
Chapter 4	Data Engineering	107
4.1	Dataset Compilation	108
4.2	Physics of Data	109
4.3	Four Pillars Framework	111
4.4	Data Acquisition	118
4.5	Data Pipeline Architecture	123
4.6	Systematic Data Processing	135
4.7	Data Labeling	143
4.8	Storage Architecture	148
4.9	Operational Data Health	158
4.10	Fallacies and Pitfalls	161
4.11	Summary	162
Part II: Development and Training		
Development Principles		167
Chapter 5	Neural Computation	169
5.1	From Logic to Arithmetic	170
5.2	Computing with Patterns	172
5.3	Neural Network Fundamentals	183
5.4	Learning Process	196
5.5	Inference Pipeline	213
5.6	Handwritten Digit Recognition	216
5.7	D·A·M Taxonomy	220
5.8	Fallacies and Pitfalls	220
5.9	Summary	222
Chapter 6	Network Architectures	225
6.1	Architectural Principles	226
6.2	MLPs: Dense Pattern Processing	231
6.3	CNNs: Spatial Pattern Processing	236
6.4	RNNs: Sequential Pattern Processing	246
6.5	Attention: Dynamic Processing	250
6.6	Transformers: Parallel Sequence Processing	256
6.7	Sparse Architectures: RecSys	261
6.8	Shared Building Blocks	263
6.9	Computational Primitives	270
6.10	Architecture Selection Framework	277

6.11	Fallacies and Pitfalls	283
6.12	Summary	284
Chapter 7	ML Frameworks	287
7.1	Three Framework Problems	288
7.2	The Ladder of Abstraction	290
7.3	Execution Problem	291
7.4	Differentiation Problem	306
7.5	Abstraction Problem	317
7.6	nn.Module Abstraction	330
7.7	Framework Platform Analysis	334
7.8	Deployment Targets	339
7.9	Framework Selection	340
7.10	Anatomy of a Training Step	343
7.11	Fallacies and Pitfalls	346
7.12	Summary	348
Chapter 8	Model Training	351
8.1	Training Systems Fundamentals	352
8.2	Iron Law of Training Performance	354
8.3	Mathematical Foundations	357
8.4	Pipeline Architecture	372
8.5	Identifying Bottlenecks	382
8.6	Pipeline Optimizations	384
8.7	Scaling Training Systems	405
8.8	Fallacies and Pitfalls	412
8.9	Summary	414
Part III: Optimization and Acceleration		
Optimization Principles		419
Chapter 9	Data Selection	421
9.1	Data Selection Fundamentals	422
9.2	Static Pruning	428
9.3	Dynamic Selection	432
9.4	Self-Supervised Learning	437
9.5	Synthetic Data Generation	439
9.6	Decision Framework	442
9.7	Selection Engineering	444
9.8	Cost Modeling	447
9.9	Distributed Selection	450
9.10	Cross-Layer Interactions	452
9.11	Measurement Framework	454
9.12	Fallacies and Pitfalls	457
9.13	Summary	459
Chapter 10	Model Compression	461
10.1	Optimization Framework	462
10.2	Deployment Context	464
10.3	Structural Optimization	467
10.4	Quantization and Precision	485
10.5	Architectural Efficiency	504
10.6	Technique Selection	515
10.7	Optimization Strategies	517
10.8	Efficiency Measurement	518

10.9	Implementation Tools	520
10.10	Fallacies and Pitfalls	522
10.11	Summary	524
Chapter 11	Hardware Acceleration	527
11.1	Acceleration Fundamentals	528
11.2	Hardware Specialization	531
11.3	AI Compute Primitives	540
11.4	Compute Units and Execution Models	547
11.5	AI Memory Systems	559
11.6	Roofline Model	571
11.7	Hardware Mapping	577
11.8	Dataflow Optimization	582
11.9	Compiler Support	597
11.10	Runtime Support	603
11.11	Multi-Chip Scaling	605
11.12	Heterogeneous SoC Design	607
11.13	Hardware Sustainability	609
11.14	Fallacies and Pitfalls	610
11.15	Summary	612
Chapter 12	Benchmarking	613
12.1	ML Benchmarking Framework	614
12.2	Historical Foundations	616
12.3	System Benchmarking Suites	619
12.4	Benchmarking Granularity	625
12.5	Benchmark Components	629
12.6	Training vs. Inference	638
12.7	Training Benchmarks	639
12.8	Inference Benchmarks	645
12.9	Power Measurement Techniques	656
12.10	Benchmarking Best Practices	661
12.11	Model and Data Evaluation	666
12.12	Production Considerations	674
12.13	Fallacies and Pitfalls	676
12.14	Summary	677
Part IV: Deployment and Operations		
Deployment Principles		681
Chapter 13	Model Serving	685
13.1	Serving Paradigm	686
13.2	Serving Load, Latency, and Architecture	686
13.3	Serving System Architecture	693
13.4	Request Lifecycle	695
13.5	Queuing Theory	702
13.6	Model Lifecycle Management	707
13.7	Throughput Optimization	711
13.8	LLM Serving	722
13.9	Inference Runtime Selection	725
13.10	Node-Level Optimization	728
13.11	Economics and Planning	731
13.12	Fallacies and Pitfalls	734
13.13	Summary	736

Chapter 14 ML Operations	739
14.1 MLOps Overview	740
14.2 Principles and Foundations	741
14.3 Technical Debt	746
14.4 Development Infrastructure	751
14.5 Production Operations	763
14.6 Design and Maturity Framework	788
14.7 Case Studies	792
14.8 Fallacies and Pitfalls	798
14.9 Summary	799
Chapter 15 Responsible Engineering	803
15.1 Responsibility as Systems Engineering	804
15.2 Engineering Responsibility Gap	805
15.3 Responsible Engineering Checklist	814
15.4 Environmental and Cost Awareness	826
15.5 Data Governance and Compliance	831
15.6 Fallacies and Pitfalls	836
15.7 Summary	837

Synthesis

Chapter 16 Conclusion	841
16.1 Synthesizing ML Systems	842
16.2 Thirteen Quantitative Principles	844
16.3 Principles in Practice	849
16.4 Future Directions	850
16.5 Journey Forward	853
16.6 Fallacies and Pitfalls	854
16.7 Summary	855

Back Matter

Appendices	859
Appendix A The D·A·M Taxonomy	859
Appendix B Data Foundations	867
Appendix C Algorithm Foundations	873
Appendix D Machine Foundations	887
Appendix E System Assumptions	899
Glossary	911
References	923
Index	951

Foreword

David A. Patterson

Pardee Professor of Computer Science Emeritus
University of California, Berkeley
Distinguished Engineer, Google
ACM A.M. Turing Award Laureate
Berkeley, California

My career has been dedicated to understanding and improving the very foundations upon which computing operates – from the instruction set architectures that define processors to the storage systems that hold our data. I’ve seen firsthand how profound advancements are rarely about a single brilliant idea in isolation, but rather about the holistic integration of many ideas into a robust, efficient, and reliable system.

Today, we stand at another pivotal moment. Machine learning (ML) has transcended scientific curiosity to become a force reshaping industries, driving innovation, and transforming our daily lives. The intellectual power of the algorithms, the gains in accuracy, and the seemingly magical capabilities they unlock have rightly captured the world’s imagination.

One concern is that much of this progress has been achieved by a community that, through no fault of its own, was never taught to think in systems terms about what it was building. Engineers who can call `model.fit()` but cannot explain why training is slow. Researchers who tune hyperparameters without understanding the memory bandwidth constraints that make certain choices impractical. Practitioners who deploy models into production without a framework for reasoning about latency, throughput, or failure modes.

A brilliant algorithm, without a well-engineered system to support it, is like a powerful engine without a chassis and brakes. It won’t get you far, and it certainly won’t get you there reliably.

This gap is not a talent deficit. It is a curriculum deficit. The field has simply not had the textbooks it deserves. This book, *Introduction to Machine Learning Systems*, arrives precisely at the right time to illuminate this critical, often overlooked, dimension of the ML revolution.

Why Systems Thinking Matters

My career has leveraged the power of systems thinking. Identify the bottlenecks, understand the interactions between components, and design for performance, reliability, cost-effectiveness, and ease of use across the entire stack. You could not reach the conclusion from first principles alone. You had to think about the system.

ML systems demand exactly the same discipline. A neural network is not a mathematical abstraction floating free of physical constraints. It is a computation that must execute on silicon with finite memory bandwidth, finite power budgets, and finite communication capacity. Arithmetic intensity, the Roofline Model, and memory hierarchy effects are not advanced topics to be addressed after students learn the basics. They are the basics.

Why ML Systems Deserve Serious Textbook Treatment

The purpose of a textbook is not just to catalogue the current frontier. It is to provide the conceptual scaffolding that allows students to understand the frontier, and to advance it themselves. The neural network architectures covered in these pages will evolve. The hardware accelerators described will

be superseded. But some fundamentals will not change. We will still have memory hierarchies and arithmetic intensity will still be a critical measure. The principles of pipelining, of latency versus throughput trade-offs, of quantization and precision reflect physical and mathematical realities that do not go out of date. Anyone who genuinely understands them will be able to reason about whatever new system the field produces next year or ten years from now.

I am also struck by the scope of what this book attempts. It is not enough to explain how training works; one must also address data engineering, model optimization, hardware acceleration, deployment, benchmarking, responsible AI, and sustainability. These are not separate topics bolted together for completeness. They are interconnected aspects of a single engineering discipline. A decision made at the data layer has implications for model behavior, which has implications for inference latency, which has implications for deployment cost, which has implications for the communities that can and cannot access the resulting system. The book correctly treats this integration as a first-class concern.

I am particularly pleased to see the quantitative methodology that runs through every chapter. Claims are backed by measurements. Trade-offs are expressed numerically. Students are taught to estimate, measure, and verify rather than to rely on their gut. The difference between an engineer and a dilettante is precisely this capacity for rigorous, measurement-grounded reasoning.

Why Open Access Is Important for the Next Generation

I applaud the decision to make this invaluable resource freely available.

When I co-wrote *Computer Architecture: A Quantitative Approach* with John Hennessy, we debated how to make the material as accessible as possible. We lobbied our publisher to keep the price down and even came out with a low cost international student edition for sales outside North America. Those concerns have only grown more urgent as ML has become infrastructure — as consequential to the modern world as the electrical grid or the internet.

The decision to publish this book as a freely available educational resource is not a compromise. It is a statement of values. ML systems are being deployed in every sector of every economy on every continent. The engineers who build and maintain them will come from everywhere. If the foundational education required to do this work responsibly is available only to students at well-resourced institutions, we will contribute to an expertise gap that mirrors the societal concerns of who benefits from AI and who bears its costs.

What impresses me further is that this project has gone well beyond open access to a static text. The accompanying ecosystem — interactive laboratories, a build-from-scratch framework in TinyTorch, hardware deployment kits, instructor materials, a global network of adopting universities — reflects a familiar educational maxim: “I hear and I forget. I see and I remember. I do and I understand.”

The open-source model also makes the curriculum itself better. Errors get corrected. Examples get sharpened. New chapters emerge in response to what the field actually needs. As we found with *Computer Architecture: A Quantitative Approach*, a community of contributors that forms around a book can offer a form of quality assurance that no single author team, however expert, can match alone.

A Final Word

The field of computer architecture has a tradition I am proud of: we build things, measure them honestly, share what we learn, and revise our philosophies when the measurements demand it. We do not hide behind the complexity of our systems to avoid accountability for whether they actually work. We do not conflate novelty with progress. We debate, rigorously, about what matters and why.

Hopefully, ML systems engineering will follow this tradition. It needs the habits of mind that come from taking systems seriously: thinking about interactions, measuring at the right level of abstraction, being honest about trade-offs, and holding oneself accountable to real-world performance rather than benchmark folklore. This book embodies those habits.

Introduction to Machine Learning Systems is more than just a textbook; it is a clarion call for systems thinking in the age of AI, a guide for building robust ML applications, and a testament to the power of open education.

We are at an extraordinary moment. Over my career there has been a major breakthrough every decade: the microprocessor in the 1970s, the personal computer in the 1980s, the Internet in the 1990s, the smartphone in the 2000s, and AI in the 2010s. The first four changed business and society, yet AI looks likely to have the largest impact of them all.

The systems being built today will shape how intelligence is deployed, who benefits from it, and what it costs — in money, in energy, and in carbon footprint. We need the engineers who build those systems to have the best possible systems skills.

This book is a significant step toward making it possible for anyone to learn them.

David A. Patterson

Author's Note

THE WORLD IS RUSHING to build AI systems; it is not yet engineering them. That sentence is the reason this book exists. A model that wins on a benchmark is not yet a system, and the distance between the two is not a detail of implementation. It is engineering: the work of making something hold up under constraints that do not negotiate. A demo answers to no one. A system answers to latency budgets, memory limits, power envelopes, and the cost of being wrong in front of real users.

What makes this its own discipline is an unusual coupling. A processor answers to the physics of the machine, to clock cycles, cache lines, and the heat it can shed. A statistical model answers to the mathematics of learning, to bias, variance, and how generalization changes with data and model capacity. A machine learning system must answer to both at once, in the same decision, and the two do not always agree. A change that pleases the hardware can starve the model of the precision or throughput it needs to converge; a change that helps the model can ask the silicon for more memory or bandwidth than it can give. Hold that tension in view and the field stops looking like a pile of tools and starts looking like a structure you can reason about.

This book hands you that structure. Not the framework that is fashionable this year or the accelerator that is fastest this quarter, but the way of thinking that outlasts them. Frameworks change, hardware turns over every few years, and the benchmarks reset; the forces underneath stay the same. If the book works, you will finish it able to look at any machine learning system, from a sensor sipping microwatts to a server farm drawing megawatts, and see the same small set of forces deciding what is possible.

The textbooks that changed me when I was a student did not hand me facts. They handed me a way of seeing, and the facts arranged themselves around it. That is what I have tried to do here.

— Vijay Janapa Reddi
Cambridge, Massachusetts
2026

Preface

Who This Book Is For

This book is for anyone who wants to understand *how* machine learning systems actually work, from the mathematics that define a neural network to the hardware that executes it and the engineering decisions that make it run efficiently.

Most readers will not become ML hardware architects or compiler engineers. Yet every practitioner who builds, trains, optimizes, or deploys a model makes decisions that depend on understanding the system beneath the code. This book provides that systems foundation for undergraduates encountering ML systems for the first time, practicing engineers strengthening their foundations, and researchers who need to reason about performance and deployment.

Why a Systems Textbook

In 1968, a NATO conference in Garmisch, Germany, helped establish *software engineering* as a distinct discipline amid concern that software systems had become too complex for ad hoc programming (Naur and Randell 1969). Reliable software required its own principles, methods, and rigor. Decades later, Hennessy and Patterson made quantitative analysis central to computer architecture education, systematizing how engineers measure, predict, and optimize processor performance (Hennessy and Patterson 2017). Both examples show how shared methods make complex engineering work systematic and teachable.

Machine learning faces its own Garmisch moment.

For the past decade, ML research has been dominated by *model-centric* thinking focused on discovering new algorithms and architectures. That focus matters, but a model is not a system. A neural network that performs well in a research notebook is useless if it can not meet latency constraints, if its data pipeline can not handle the required volume, or if its energy cost exceeds its economic value. The gap between a working model and a production system is not merely an implementation gap; it is a gap in fundamental engineering principles.

This book treats ML systems engineering not as a collection of tools, such as Kubernetes, PyTorch, and CUDA, but as a discipline grounded in quantitative bounds, models, and design principles. Just as civil engineers rely on foundational physics and structural principles, AI engineers must make the governing principles behind each system model explicit:

- **The iron law of ML systems:** A decomposition of movement, compute, and overhead, with overlap handled by the critical path.
- **Data gravity:** The placement pressure created by large data volumes and constrained networks.
- **Energy accounting:** Per-event costs and event counts determine whether movement or computation dominates.
- **Amdahl's law:** The serial fraction of a pipeline caps total system speedup.

Computer science studies what is computable. ML research studies what is learnable. AI engineering studies what is buildable, given the physics of real hardware and real data.

I wrote this book to help build that foundation and to do for AI systems what Hennessy and Patterson (2017) did for computer architecture: replace intuition with measurement, and art with engineering. The goal is to teach system-level reasoning before implementation, treating data, algorithms, and hardware not as separate silos but as a single, coupled optimization problem.

Consider a box of LEGO bricks. The same interlocking pieces build a spaceship, a castle, or an entire city. The sets change; the bricks do not. ML systems work the same way. A convolutional network for image classification, a transformer for language generation, and a reinforcement learning agent for robotics are vastly different applications, but beneath them sit the same building blocks: computational graphs, memory hierarchies, data pipelines, optimization loops, and the trade-offs between latency, throughput, and energy. Mastery of those building blocks makes unfamiliar systems legible.

That is why this book teaches enduring principles rather than current tools. Dominant frameworks have shifted from Theano to TensorFlow to PyTorch in under a decade. Hardware has evolved from repurposed graphics processors to purpose-built tensor accelerators, with new architectures emerging every year. Any textbook that teaches today's tools will be obsolete before its next edition. The building blocks endure.

What This Book Covers


This book centers on the *single compute node*: one host with one to eight accelerators, typically with discrete device memories connected by an interconnect. It also introduces multi-accelerator training and production operations where they clarify node-level design; fleet-scale treatment lies outside its scope.

The content is organized into four parts, each with a distinct role in the systems argument:


- **Part I: Foundations of ML Systems** develops the vocabulary, mental models, and end-to-end workflow that guide every ML project, from the characteristics that distinguish ML systems from traditional software through the data engineering pipelines that feed them.
- **Part II: Development and Training** connects mathematical foundations to working systems through deep learning computation, neural architectures, frameworks, and training procedures.
- **Part III: Optimization and Acceleration** addresses data efficiency, model compression, hardware acceleration, and the benchmarking methodology needed to measure progress rigorously.
- **Part IV: Deployment and Operations** covers serving infrastructure, operational practices, and responsible engineering for systems that must be fair, transparent, and trustworthy in production.

These four parts trace a path through the **ML Systems Stack**, a layered abstraction modeled on the classical computer systems stack. Each layer provides an abstraction to the layer above and consumes the layer below: from hardware at the bottom through frameworks, models, training, and serving, up to operations and applications at the top. Data flows through all layers as a cross-cutting interconnect, much like a data bus in computer architecture.

Throughout the book, a margin figure highlights which layers each chapter addresses, providing a recurring map of the stack.




The data bar alongside the stack is not merely decorative. Its uniform shading reflects how central data is to each chapter's concerns, a separate dimension from the layer intensities. The connecting wires between each layer and the data bar show which layers interact with data in that chapter's context. Three chapters illustrate how the lens shifts.



In chapter 11, the layer intensity concentrates at the bottom of the stack while the data bar is faint: the chapter is about silicon, not data. In chapter 8, the middle layers glow and the data bar is moderately lit, reflecting that training consumes data but is fundamentally about optimization. In chapter 15, the top layers dominate and the data bar carries a moderate tint, because biased training data surfaces as biased applications even though several layers sit between them. The same stack tells three different stories, with data as an independent dimension.

The margins also carry small visual notes at the point where a relationship is easiest to understand by sight. These are not numbered figures to cite later; they are reading aids. A ladder makes a scale gap visible, a trend sketch shows direction, a roofline thumbnail locates a bottleneck regime, and a dominance bar shows which term controls a total. When a miniature uses a log scale, it says so inside the image.



Each part builds on the previous one. Sequential reading is recommended for those new to the field, while the following tailored paths support specific professional goals:

Suggested Reading Paths

Readers with different backgrounds and goals may benefit from tailored paths:

- **The Complete Path** (all chapters, in order): For undergraduates and readers new to ML systems. Start with Part I, then progress through Parts II, III, and IV. This path develops every concept from first principles.

- **The ML Engineer’s Path** (systems depth for practitioners): For engineers who already train and deploy models but want to understand the systems underneath. Skim Part I for vocabulary, then focus on chapter 7, chapter 8, chapter 11, chapter 12, and chapter 13.
- **The Optimization Path** (efficiency-focused): For engineers working on model efficiency, edge deployment, or resource-constrained systems. After Part I, move directly to chapter 9, chapter 10, chapter 11, and chapter 12.

Conventions

This book uses a small set of typographic conventions to mark what each piece of prose is doing.

Bold marks first definitions

A term in **bold** within body prose is its formal first introduction in the volume. Each term receives this treatment exactly once. Subsequent appearances are unbolded; by then the term is part of the working vocabulary the chapter carries forward.

Every machine learning system has a **dual mandate**: it must manage statistical uncertainty and physical constraints simultaneously.

The bold is a contract: register this term because it will recur. A matching index entry in the back of the book points to the defining occurrence (under the **definition** subentry), providing a way to return to it. Key technical terms are also compiled in the glossary.

Bold also appears in a second, structural role: as a functional label at the start of an item inside a callout box or a summary bullet—for example, **Recommendation**, **Trade-off**, or **Result**. In that context, the bold is scaffolding rather than a first definition.

Italics carry specific rhetorical jobs

Italics are not used for general emphasis. Every italic span in this book performs one of a small number of specific jobs:

- **Direct opposition**: *training vs. inference, logical vs. physical.*
- **Impossibility, stated as physical, mathematical, or logical law**: “sub-10 ms systems *cannot* use distant cloud infrastructure—physics forbids it.”
- **Word used as word**: “the term *gradient* refers to a vector of partial derivatives.”
- **Titles of works**: *Computer Architecture: A Quantitative Approach.*
- **Thesis punchline**: A single italicized sentence that captures the controlling insight of a section, for example *Data is Source Code.*

An italicized word is performing one of these jobs. A word that is not italicized carries no implicit emphasis.

Callouts have visual identities that match their job

Callout boxes are not decoration. Each type signals a specific kind of content, so the visual identity announces what the box contains before the reader engages with it.

Callout	Contains
Definition	Formal definition of a key term, with its significance, distinction from the nearest neighbor concept, and a common pitfall
Example	Historical case or concrete scenario, structured as context, insight, and systems lesson
Systems Perspective	Analytical observation or mathematical nuance that enriches the surrounding argument
Napkin Math	Worked calculation walked through step by step, with values computed inline from the book’s constants; intermediate tables and equations are the math itself
Checkpoint	Self-check questions to verify understanding of the preceding section

Callout	Contains
War Story	Documented engineering incident with a sourced failure mode, structured as context, failure, consequence, and lesson
Lighthouse	Recurring workloads characterized by a profile (parameters, FLOPs, dominant constraint), used as anchored case studies across multiple chapters
Principle	Canonical book principle or invariant referenced from later chapters
Theorem	Formal, equation-backed result or bound, stated with its conditions and applied in later quantitative analysis
Key Takeaways	Four to seven most important results from a chapter, in summary form
Learning Objectives	Skills and concepts that readers should be able to explain or apply after completing the chapter
What's Next	Bridge from the chapter just ended to the chapter beginning, naming the open question the next chapter addresses

When a callout contains a table, figure, or equation, that artifact is part of the callout's pedagogical unit, not decoration. A lighthouse's profile *is* its characterization. A napkin math box's intermediate tables *are* the worked math. Reading these as separated artifacts loses the teaching arc. Standalone reference tables, by contrast, sit between paragraphs as numbered artifacts the reader returns to through cross-references. Both forms are valid; the visual marker (callout vs. standalone) tells the reader which one they are looking at.

Whichever path readers follow, the text is meant to connect with supporting resources beyond the chapters themselves.

Companion Resources

This book is part of a broader learning ecosystem. The complete text, interactive notebooks, lecture slides, exercises, hands-on labs, and supplementary materials for every chapter are available online at mlsysbook.ai.

This book is open source. The full text, figures, and build system are publicly available, and contributions from readers worldwide have materially improved every chapter. Visit mlsysbook.ai/git to report issues, suggest improvements, or contribute directly.

Prerequisites

Readers should bring two prerequisites:

- **Programming:** Fluency in Python, including functions, classes, and basic data manipulation with NumPy.
- **Mathematics:** Comfort with linear algebra, basic calculus, and probability at the undergraduate level.

Two additional areas are helpful but not required:

- **Computer systems:** Familiarity with memory hierarchies and basic computer architecture deepens the optimization and hardware chapters.
- **Machine learning basics:** Prior exposure to supervised learning concepts provides useful context, but Part II develops the necessary foundations from first principles.

Beyond This Book

This book develops the single-node foundation while introducing distributed execution, production operations, and governance where they affect the end-to-end system. Advanced multi-node topologies, fleet-wide collective algorithms, and cluster resilience extend into distributed systems at scale. Completing this text prepares readers to navigate that broader engineering landscape.

Using This Book in a Course

This book grew out of CS249r at Harvard University and the TinyML edX professional certificate program. The TinyML course taught ML systems at the most constrained end of the spectrum:

microcontrollers running on milliwatts with kilobytes of memory. It revealed a surprising lesson. The fundamental constraints that govern tiny devices (memory bandwidth, compute density, energy per operation) are the same constraints that govern data center GPUs. *The physics scales; only the numbers change.* That insight shaped this book. What began as a course on resource-constrained ML matured into a comprehensive treatment of ML systems engineering, grounded in the principle that mastery of the fundamentals at any scale prepares engineers for every scale.

In course settings, the book is designed to support a one-semester sequence covering the full ML systems stack. Instructors may also select individual parts for shorter modules:

- **Parts I–II** form an introductory half-semester on ML systems foundations and how such systems are built and trained.
- **Parts III–IV** form an applied half-semester on optimizing, deploying, and operating them.

For survey courses or executive programs, the four part introductions alone provide a condensed overview of the quantitative principles that govern ML systems.

Lecture slides, assignments, and instructor resources are available at mlsysbook.ai.

Acknowledgments

Origins

This book owes its existence to a single conversation. During a sabbatical at Google, Pete Warden invited me to work on what seemed like an impossible moonshot for TinyML: running a small speech model on a microcontroller. That collaboration on TensorFlow Lite for Microcontrollers revealed something fundamental. The constraints that govern a device with 16 KB to 256 KB of memory are the same constraints that govern an accelerator in a data center, even when the scale is entirely different. Bandwidth, latency, energy, and memory obey the same physics across both settings. That insight became the seed of this book.

The TinyML work led to an online course on HarvardX, co-taught with Laurence Moroney, that grew to over 100,000 students worldwide. Brian Plancher helped build the TinyML curriculum. Marcelo Rovai, a professor who first encountered the material through the TinyML course, went on to develop the hands-on labs and kits that accompany this book. Zach Shelby opened the Edge Impulse platform for student use. Massimo Banzì and the Arduino team, along with Eric Pan and Seeed Studio, collaborated with us in developing and providing the hardware kits that made hands-on instruction possible. Marco Zennaro brought the material to developing countries through ICTP and planted the idea that these principles deserved a proper textbook. Teaching at that scale, without one, made the need undeniable.

That need turned into a book through CS249r at Harvard University. In 2023, the semester-long class project was to write the class notes: each student group took ownership of a topic, researched it, and worked with me to develop it into a chapter. That work seeded every chapter in this book. The premise was that writing forces you to think deeper than reading ever does. We then opened it all on GitHub, and what began as rough class notes has been rewritten and refined with every semester since.

Because that GitHub work continues, this book benefits from a community whose contributions evolve with the project. New contributors join as readers report issues, improve explanations, and extend the supporting materials. A printed acknowledgment can capture only one moment in that continuing history. The current contributor list and contribution guidelines are available at <https://mlsysbook.ai/git>.

Collaborators and Colleagues

Beyond the class and GitHub community, my earlier collaborations shaped the book's intellectual foundations. My time at Google informed much of the systems thinking in these pages. That period began with performance modeling for the first Pixel system on chip (SoC), drawing on my background in computer architecture. A collaboration with Mark Hill on roofline-style analytical modeling for heterogeneous SoCs taught me to reason about hardware from first principles, the kind of quantitative thinking that runs through every chapter. That emphasis on analytical modeling led naturally to systematic evaluation, and systematic evaluation led to MLPerf. MLPerf involved many contributors across MLCommons, but I worked most closely with Peter Mattson, David Kanter, Carole-Jean Wu, Christine Cheng, Guenther Schmuelling, Cody Coleman, and Cliff Young. That collaboration shaped how I think about rigorous measurement of ML systems.

Boris Murmann and Song Han convinced me in the early stages that this project was worth pursuing. As I wrote, I sought feedback from colleagues in industry and academia alike. Practitioners at Arm, Google, Intel, Meta, Microsoft, Qualcomm, STMicroelectronics, and other organizations

offered perspectives grounded in production systems, while academic reviewers at universities across Europe, Asia, and the Americas stress-tested the material against their own research and teaching. Their combined feedback sharpened every chapter. Evgeni Gousev at Qualcomm and Pete Bernard at the Edge AI Foundation supported the project's outreach and helped build the community around it.

That feedback loop extended to classrooms around the world, where instructors put the material to the test when it was still rough. Sophia Shao at UC Berkeley, Xiaofan (Fred) Jiang at Columbia, Tushar Krishna and Divya Mahajan at Georgia Tech, and Wenkai Guan at the University of Minnesota Morris were among the first, and instructors at IIT Bombay, NUST Pakistan, De La Salle University in the Philippines, and universities across Europe, Latin America, and Africa have since adopted the book for their own courses. Their willingness to teach from an unfinished textbook provided invaluable feedback and confirmed that the need for this material was not limited to any one institution or any one country.

Students and Contributors

My students in the Edge Computing Lab at Harvard shaped this book in ways I can not fully trace. Some did foundational work that helped establish TinyML as a field; their research and the ideas that emerged from years of collaboration informed the book's technical content. This cohort helped set the field in motion and then helped me write about it: Andy Cheng, Arya Tschand, Brian Plancher, Colby Banbury, Elizabeth Suitor, Emma Chen, Ikechukwu Uchendu, Jason Jabbour, Jason Yik, Jeffrey Ma, Jessica Quaye, Mark Mazumder, Matthew Stewart, Maximilian Lam, Radhika Ghosal, Shvetank Prakash, Srivatsan Krishnan, Yu-Shun Hsiao, and Zishen Wan. Kai Kleinbard, then at the Harvard Extension School, built the online generative AI SocratiQ tutor bot. Naeem Khoshnevis, from the Harvard Kempner Institute, established essential infrastructure including GitHub Actions and documentation standards. Željko Hrček produced the TikZ diagrams and contributed to the PDF layout, page composition, and final visual formatting of the book.

A textbook that distills principles from an active research field also owes a debt to the researchers whose work it builds on. The open-source community shaped this book in ways that no single author could. Dozens of contributors caught errors, improved explanations, filed issues, and submitted pull requests. Some fixed a single typo; others rewrote entire sections. The book improved because people who owed it nothing chose to make it better. I am grateful to every one of them.

Support

This work also depended on support from the Harvard Data Science Initiative, Harvard Extension School, and the National Science Foundation; from the Edge AI Foundation and ICTP for educational outreach and scholarships; and from Arduino, Edge Impulse, Google, and Sseed Studio for hardware, tooling, and infrastructure that enabled the practical labs.

At MIT Press, Susan Hartman believed in this project and guided it from an open-source experiment to a published textbook. Her support, editorial judgment, and friendship made this book real in a way that GitHub alone could not. I am equally grateful to the MIT Press for agreeing to keep this book available as open access. A textbook about a discipline shaped by the open-source community should be accessible to everyone in it. That MIT Press shared this conviction made all the difference.

Finally, I thank my wife Kari and my daughters Nora and Maya. This book was written in the hours that belonged to them: late nights that started at ten and ended at two or three in the morning, weekends that disappeared into revisions, years of mornings when I was present but not quite there. What kept me going, aside from stubbornness, was the open-source community. Every GitHub star was a signal that someone out there cared, that someone was learning, that the work mattered. The cost was real, and they paid it with patience I did not always deserve.

A Note on AI Assistance

Every chapter of this book was researched, written, and rewritten by me. I also used AI tools throughout. What follows is more than the disclosure MIT Press requires of authors who use AI. It is a statement of the standard I held the book to, and a claim about what authorship means when the tools are this capable.

The standard is traceability. Derived quantities in these chapters are backed by Python computation cells whose source code ships with the book. Citations were checked against their sources, and cross-references are checked against stable anchors. The purpose-built infrastructure modeling engine used for the book's system calculations is available at <https://mlsysbook.ai/mlsysim>. A pre-commit test suite checks the machine-testable parts of these invariants on every commit, covering unit consistency, inline-reference resolution, cross-reference integrity, notation canonicity, and citation hygiene. AI tools made parts of this rigor mechanically tractable. The enforcement infrastructure makes that work auditable.

Authorship is thinking. Which concepts belong in this book and which do not. How to sequence the chapters so that each idea arrives only after the reader has the tools to understand it. Which worked examples to trace end-to-end, and which numbers to compute so the reader can verify the argument independently. What level to pitch each explanation at, rigorous enough for a graduate engineer and concrete enough for a student encountering systems thinking for the first time. Where students will get stuck, because I have taught this material and watched them get stuck. These are judgment calls that no tool can make, because they require knowing the reader. I used AI tools throughout this process to brainstorm framings, explore alternatives, and pressure-test my reasoning. The choices are mine.

In practice, AI was genuinely useful for writing code: the computation cells, the pre-commit checks, and the worked examples that are essentially programs. It helped me survey literature I then read in the original, draft passages I then rewrote substantially, and audit the manuscript for consistency across hundreds of cross-references, thousands of index entries, and tens of thousands of lines of prose. Throughout, the tools proposed and I ratified.

This pattern is older than the present moment. In 1683 Joseph Moxon printed the first English-language manual on printing, *Mechanick Exercises on the Whole Art of Printing*, using the press it described. Three centuries later Donald Knuth wrote *The TeXbook* in TeX, the typesetting system he created after becoming dissatisfied with the typesetting of his own books. A book about ML systems infrastructure, written without the infrastructure it describes, would be a contradiction. The toolmaker documents the tool with the tool.

Machine learning systems can not be certified the way traditional logic can. That verification gap applies to AI-assisted authoring with equal force. The standard I held this book to is the standard I want us to demand of every system we ship: the tool can propose, but only an engineer can ratify. MIT Press does not list AI tools as authors because they can not assume ethical and legal responsibility for their work. The same is true of every engineered system. Authorship is the assumption of that responsibility, and it belongs to the human.

Notation

Machine Learning Systems spans *machine learning* (computer science and statistics) and *systems* (computer architecture and hardware). Each field developed its notation independently, and many symbols mean different things across communities and publications. This collision creates real confusion when the disciplines merge. The conventions below establish a single notation that eliminates this ambiguity.

Consider a simple statement: “Increasing B improves throughput.” To an ML researcher, B means batch size. To a hardware engineer, B means bandwidth. Both interpretations are correct in their respective fields, but in ML Systems we need both concepts in the same equation—hence the need for a single, consistent convention.

The Iron Law of ML Systems

For serialized execution phases, the iron-law performance decomposition is:

$$T = \underbrace{\frac{D_{\text{vol}}}{\text{BW}}}_{\text{Memory Time}} + \underbrace{\frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}}_{\text{Compute Time}} + \underbrace{L_{\text{lat}}}_{\text{Latency Overhead}}$$

Each variable was chosen deliberately to avoid collision with standard ML terminology.

Symbol	Definition	Unit	Why This Symbol?
T	Time	seconds	Unambiguous. Wall-clock time for an operation.
D_{vol}	Data Volume	bytes	Avoids collision with D (Dataset Size). In scaling laws, D means training tokens. Here we need bytes moved through memory. The subscript disambiguates.
BW	Bandwidth	bytes/s	Avoids collision with B (Batch Size). Systems literature often uses B for bandwidth, while ML literature often uses B for batch size. The notation preserves the ML convention.
O	Operations	FLOPs	Total floating-point operations. Clean in equations (vs. “Ops”).
R_{peak}	Peak Rate	FLOP/s	Avoids collision with P (Parameters). Roofline presentations often use P for performance, while ML literature commonly uses P for parameter count. The notation preserves the ML convention.
η_{hw}	Efficiency	—	Hardware utilization ($0 \leq \eta_{\text{hw}} \leq 1$). Avoids collision with learning rate (η).
L_{lat}	Latency	seconds	Avoids collision with $\mathcal{L}$ (Loss). Fixed overhead time (kernel launch, network RTT). The subscript distinguishes from the loss function.

Why these choices matter

Without careful notation, sentences become ambiguous:

“Reducing D improves performance.”

The sentence has two possible readings:

- Reducing **dataset size** (fewer training samples) can speed training but may reduce accuracy.
- Reducing **data volume moved** (for example, through compression or quantization) can speed inference, with accuracy effects that depend on the technique and calibration.

With our notation, we can write precisely:

“Reducing D_{vol} through FP32-to-INT8 quantization cuts parameter memory traffic to one quarter while D (training data) remains unchanged.”

Our notation eliminates this ambiguity: “*BW limits throughput*” has only one reading.

Subscripted variants

An upright root names the quantity, and the subscript names which instance of it is meant. The convention holds closely related measurements apart instead of collapsing them into a single symbol:

- $BW_{\text{disk}}, BW_{\text{network}}, BW_{\text{accelerator}}$: bandwidth at three points on the data path
- $D_{\text{vol}}, R_{\text{peak}}, L_{\text{lat}}, \eta_{\text{hw}}$: iron-law terms, each held clear of a common ML symbol
- $E_{\text{move}}, E_{\text{compute}}, E_{\text{total}}$: one energy budget split into tradeable parts

The subscript therefore carries part of the claim. Disk and accelerator bandwidth differ by orders of magnitude, so a figure is reproducible only when the subscript says where it was measured.

The Degradation Equation

The degradation equation is a fitted local model. Divergence alone does not determine the direction of accuracy change, so estimate λ from labeled outcomes. Some symbols below, such as τ , appear in chapter prose rather than in the equation.

$$\text{Accuracy}(t) \approx \text{Accuracy}_0 - \lambda \cdot \mathcal{D}(P_t \| P_0)$$

Symbol	Definition	Unit/Type	Notes
$\text{Accuracy}(t)$	Accuracy at Time t	Scalar	Model accuracy after the model has been deployed for time t .
Accuracy_0	Initial Accuracy	Scalar	Model accuracy at deployment time.
λ	Fitted Sensitivity	Scalar	Local coefficient estimated from outcomes over a stated range; not a universal constant. (<i>Not wavelength.</i>)
P_t	Current Distribution	Distribution	The data distribution at time t . (<i>Not parameters—use P for parameter count.</i>)
P_0	Training Distribution	Distribution	The data distribution at training time.
$\mathcal{D}(P_t \ P_0)$	Statistical Divergence	Scalar ≥ 0	Measures how far P_t has drifted from P_0 . Common choices: KL divergence, total variation, Wasserstein. (<i>Calligraphic to avoid collision with D = dataset size.</i>)
τ	Response Threshold	Scalar > 0	Operational threshold for investigation; it does not automatically trigger retraining.

The Energy Corollary

With effective per-byte and per-operation costs, workload energy is approximated as:

$$E_{\text{total}} \approx D_{\text{vol}} \times E_{\text{move}} + O \times E_{\text{compute}}$$

Symbol	Definition	Unit	Notes
E_{total}	Total Energy	joules	Total energy consumed by an ML workload, decomposed into data-movement and compute terms.
E_{move}	Energy per Byte Moved	joules/byte	Effective movement cost; total energy also depends on bytes moved and operations executed.
E_{compute}	Energy per Operation	joules/FLOP	Energy cost of a single arithmetic operation.

Deep Learning Notation

This book follows standard deep learning conventions (Goodfellow et al. 2016) with explicit disambiguation for systems variables.

Symbol	Definition	Dimensions/Type
B	Batch Size	Integer. The number of samples processed in parallel. (<i>Never bandwidth.</i>)
P	Parameters	Integer. The total count of scalar model parameters, including weights and biases. (<i>Never peak FLOP/s.</i>)
D	Dataset Size	Integer. Number of training samples or tokens. (<i>Never data volume in bytes—use D_{vol}.</i>)
S	Sequence Length	Integer. Number of tokens or time steps.
d	Hidden Dimension	Integer. Size of the hidden state vector.
d_{head}	Attention Head Dimension	Integer. Per-head hidden dimension in attention layers.
N_L	Number of Layers	Integer. Total number of layers in a network.
N_{heads}	Number of Attention Heads	Integer. Number of attention heads in a multi-head attention layer.
H_{KV}	Number of Key-Value Heads	Integer. Number of key-value heads in grouped-query or multi-query attention.
ℓ	Layer Index	Integer. Index for a layer. Use instead of bare L when indexing layers, since L collides with loss and latency.
$\mathcal{L}$	Loss Function	Scalar. The objective function minimized during training.
η	Learning Rate	Scalar. Step size for the optimizer. (<i>Never bare for hardware efficiency—use η_{hw}.</i>)
θ	Model Parameters	Parameter collection, often treated as a vector. The set of all learnable parameters.
$p(x)$	Distribution	Probability mass/density for a random variable. Use lowercase $p(\cdot)$ for generic distributions to avoid collision with P (parameter count).
$p(y x)$	Conditional Distribution	Conditional probability/density. Use this form for generic label relationships; reserve P_0 and P_t for the degradation equation’s training/current distributions.
$\Pr(E)$	Event Probability	Probability of an event E . Use for event statements such as $\Pr(\text{batch} = 0)$; use $p(x)$ and $p(y x)$ for distributions.

Latin matrices and vectors are set in bold ($\mathbf{W}$, $\mathbf{x}$); scalars, dimensions, indices, and individual matrix entries stay italic (N , d , i , W_{ij}). The generic parameter symbol θ is the one conventional exception and stays italic. Local matrix algebra follows the same bold rule: matrix operands are bold ($\mathbf{A}$, $\mathbf{B}$, $\mathbf{C}$), as in general matrix multiply (GEMM) $\mathbf{C} = \alpha\mathbf{AB} + \beta\mathbf{C}$. The bold form marks a matrix operand and is distinct from the italic scalar batch size B , which remains batch size in scalar contexts.

Performance, Serving, and Memory Notation

Reusable performance and serving quantities follow the same collision-avoidance rule as the iron law. Rates use R or descriptive Greek symbols; request counts use Q_{req} rather than overloading N ; queue utilization uses a subscripted ρ so bare ρ remains available for other ratio models.

Symbol	Definition	Unit/Type	Notes
I	Arithmetic Intensity	FLOP/byte	Workload FLOPs per byte moved. The Roofline Model uses I as the independent variable.
I_{ridge}	Roofline Ridge Point	FLOP/byte	$I_{\text{ridge}} = R_{\text{peak}}/\text{BW}$. Prefer this explicit form over starred shorthand so the meaning remains clear in prose.
R_{attain}	Attainable Compute Rate	FLOP/s	Roofline bound: $R_{\text{attain}} \leq \min(R_{\text{peak}}, I \times \text{BW})$. Uses R , not T , because this quantity is a rate.
MFU	Model FLOPs Utilization	Dimensionless	Useful model FLOP/s divided by available peak FLOP/s. Text acronym avoids overloading η .
r_{comp}	Compression Ratio	Dimensionless	Uncompressed size divided by compressed size. A compressed payload has $1/r_{\text{comp}}$ of the uncompressed size; the subscript avoids bare C collisions.

Symbol	Definition	Unit/Type	Notes
Q_{req}	Request Concurrency	requests	Average in-flight requests in a stable serving system ($Q_{\text{req}} = \lambda_{\text{arr}} \cdot T_{\text{lat}}$ via Little's Law). Avoids collision with N as device count in distributed settings.
λ_{arr}	Arrival Rate	requests/s	Request arrival rate. The subscript avoids collision with λ as degradation sensitivity or failure rate.
T_{lat}	Request Time in System	seconds	End-to-end queueing/serving latency for Little's Law. Distinct from L_{lat} , the fixed-latency term in the iron law.
$T_{\text{svc}}(B)$	Batch Service Time	seconds	Time to serve a batch of size B .
$\mu_{\text{eff}}(B)$	Effective Service Rate	requests/s	Batched service rate, typically $\mu_{\text{eff}}(B) = B/T_{\text{svc}}(B)$.
ρ_{serv}	Serving Utilization	Dimensionless	Queue/server utilization. Use instead of bare ρ , which is reserved for communication-computation ratio in distributed contexts.
M_{total}	Total Memory Footprint	bytes	Sum of explicit memory components; avoids bare M ambiguity.
M_{weights}	Weight Memory	bytes	Memory occupied by model parameters.
$M_{\text{gradients}}$	Gradient Memory	bytes	Memory occupied by stored gradients.
$M_{\text{optimizer}}$	Optimizer-State Memory	bytes	Momentum, variance, master weights, and related optimizer buffers.
$M_{\text{activations}}$	Activation Memory	bytes	Retained activations for the backward pass or serving intermediates.
s_{elem}	Element Storage Size	bytes/element	Bytes per stored tensor element.

Units and Precision

- **Physical units:** This book uses SI (metric) units throughout, including meters, kilograms, seconds, watts, and °C, consistent with standard engineering and scientific practice. Where source data was originally reported in imperial units, the book converts to SI and notes the original values parenthetically. A space always separates the number from the unit (for example, 100 ms, 2 TB/s).
- **Data and memory:** This book uses **decimal SI prefixes only**: KB = 10^3 bytes, MB = 10^6 , GB = 10^9 , TB = 10^{12} . Binary-prefixed units do not appear in prose; all capacities, throughputs, and model sizes are reported in decimal units (for example, 80 GB, 2 TB/s, 102 MB).
- **Compute:** The notation distinguishes operation counts from rates. Total work uses FLOPs (for example, GFLOPs, TFLOPs), while throughput uses FLOP/s with decimal prefixes (for example, GFLOP/s, TFLOP/s).
 - 1 TFLOP = 10^{12} FLOPs
 - 1 TFLOP/s = 10^{12} FLOPs per second
 - Arithmetic intensity conventionally uses FLOP/byte as a unit ratio (floating-point operations per byte moved). This is a ratio unit, not a total-work symbol or throughput symbol.
- **Currency:** Dollar amounts use the dollar sign (\$); unless otherwise noted, dollar-denominated costs are U.S. dollars (USD).
- **Precision:**
 - **FP64:** Double precision (8 bytes)
 - **FP32:** Single precision (4 bytes)
 - **TF32:** TensorFloat-32 (19-bit Tensor Core compute mode; inputs and storage remain FP32)
 - **FP16:** Half precision (2 bytes, standard range)
 - **BF16:** Brain float (2 bytes, wide dynamic range)
 - **FP8:** Quarter precision (1 byte, E4M3 or E5M2 format)
 - **FP4:** 4-bit floating-point format
 - **INT8:** 8-bit integer (1 byte)
 - **INT4:** 4-bit integer; lower integer precisions follow the same uppercase INTn pattern (for example, INT3, INT2)



Quick Reference: Resolving Collisions

Common collision points in ML Systems literature include:


Symbol	ML Meaning	Systems Meaning	Book Convention
B	Batch Size	Bandwidth	Batch Size. Use BW for bandwidth.
P	Parameters	Peak FLOP/s	Parameters. Use R_{peak} for peak rate.
D	Dataset Size	Data Volume	Dataset Size. Use D_{vol} for bytes moved.
L	Loss	Latency	Loss ($\mathcal{L}$). Use L_{lat} for latency.
η	Learning Rate	Efficiency	Learning Rate. Use η_{hw} for efficiency.

As a general principle, ML conventions take precedence for single letters, while systems concepts get subscripts or multi-letter symbols. This reflects the primary audience (ML practitioners learning systems) and preserves compatibility with the vast ML literature.

MAIN TEXT



Introduction




- 1.1 The AI Systems Moment
- 1.2 Data-Centric Paradigm Shift
- 1.3 The Evolution of AI Bottlenecks
- 1.4 The Bitter Lesson
- 1.5 Defining ML Systems
- 1.6 ML vs. Traditional Software
- 1.7 Iron Law of ML Systems
- 1.8 Three Dimensions of ML Efficiency
- 1.9 AI Engineering as a Discipline
- 1.10 Five-Pillar Framework
- 1.11 Book Organization
- 1.12 Fallacies and Pitfalls
- 1.13 Summary

Purpose

Why does building machine learning systems require engineering principles so different from those governing traditional computing systems?

Machine learning systems have a physics. Data moves through memory hierarchies governed by bandwidth, arithmetic runs on silicon governed by power, and predictions must arrive within latency windows. These constraints are not implementation details; they shape decisions from model architecture to deployment target. ML systems also differ from traditional computing systems because behavior is defined by data, not only by explicit logic or hardware state. When a conventional program misbehaves, engineers can often trace source or inspect hardware state; when an ML system misbehaves, the code may execute correctly while learned behavior fails because the data was incomplete, biased, stale, or no longer representative. ML engineers therefore manage statistical uncertainty and physical execution constraints together. A model that fits in a data center may be useless on a phone; a training pipeline that converges in a week on one accelerator may take a month on another; an accurate model trained on last year's data may silently degrade. Traditional practices such as testing, modularity, version control, and performance analysis remain necessary, but they are not sufficient. At system scale, an optimization in one layer can move the bottleneck to another, so correctness, efficiency, and deployability cannot be designed independently. The first task is therefore diagnosis, because improving the most visible component may leave end-to-end behavior unchanged when another constraint is binding. This book builds a discipline grounded in computation's physical limits. Algorithmic choices affect the stack down to the machine, and hardware constraints flow back up to model design.





Learning Objectives

- Explain why data-defined behavior and physical constraints distinguish ML systems from traditional software
- Apply a data-algorithm-machine lens to diagnose bottlenecks across data movement, arithmetic, and machine limits
- Analyze AI's shift from symbolic rules to deep learning through the bitter lesson
- Calculate iron-law performance terms to reason about throughput, latency, and return on compute
- Synthesize lifecycle, deployment, degradation, and five-pillar perspectives into ML systems engineering judgments

1.1 The AI Systems Moment

ARTIFICIAL intelligence is no longer confined to research demonstrations. Ask a smartphone a question and, within seconds, learned components convert speech to text, interpret intent, retrieve information, and generate a response. Search engines rank results, recommendation systems decide what people see, lenders use models to assess risk, and driver-assistance systems detect hazards. In each case, a prediction participates in a larger decision that can affect attention, money, access, or safety. What appears to be one intelligent action is therefore an end-to-end system operating in the world.

The modern AI movement became practical when three forces converged. Digital services generated more examples than engineers could encode as rules. Learning algorithms became able to extract useful behavior from those examples. Parallel machines made the resulting computation feasible. The most consequential change was conceptual. Data stopped being merely an input to software and became part of the mechanism that defines its behavior. Instead of writing every decision rule, engineers construct a training process through which examples shape the rules a system will apply.

This convergence creates a **dual mandate**. Every ML system must establish that its learned behavior is trustworthy and that the machine can produce that behavior within the available time, memory, energy, and cost. A model that is accurate but too slow is unusable; a fast model trained on unrepresentative data is wrong at machine speed. The distinction becomes clearest at failure boundaries. A code defect may crash loudly, while a data defect can leave every instruction executing correctly as predictions deteriorate silently. At scale, both obligations span the entire stack. Conversational services coordinate pools of GPUs¹ while managing memory, networks, and heat. Driver-assistance systems must fuse sensor streams within milliseconds. Google processes 8.5B searches per day under strict latency targets. These systems succeed only when learned behavior and physical execution are engineered together. The first part of that story is the shift from code-defined logic to data-defined behavior.

1.2 Data-Centric Paradigm Shift

That shift changes how software is built. Instead of writing behavior directly, engineers construct a process that learns behavior from data. When a traditional program fails, an engineer can often trace a branch, inspect a stack frame, and patch the code path. When an ML system loses accuracy without a code change, the cause may be a shifted data distribution, a changed label process, or a model that no longer represents production behavior. Andrej Karpathy² described this change as the move from **software 1.0** to **software 2.0** (Karpathy 2017). As table 1.1 shows, Software 1.0 encodes operational logic in instructions, while Software 2.0 learns that logic from examples. Its failures can therefore remain silent until evaluation or monitoring reveals that behavior has changed.

Software 2.0 does not eliminate code. Engineers still build data pipelines, training loops, evaluation tools, and serving infrastructure. What changes is where application behavior lives. Some of it resides in learned weights shaped by data, so code review alone cannot explain what the system will do.

¹ | **GPU (graphics processing unit):** Originally designed for rendering video game graphics, a workload requiring thousands of simple, parallel pixel calculations. This hardware-algorithm alignment proved decisive for neural networks, where the same massively parallel arithmetic structure maps directly onto matrix multiplication, making GPUs a primary physical enabler of modern training scale.

² | **Andrej Karpathy:** A founding member of OpenAI and former Director of AI at Tesla who pioneered the application of deep learning to autonomous vehicle fleets. His “Software 2.0” thesis (2017) crystallized the insight that neural network weights are the new “source code,” forcing a new engineering reality: instead of debugging explicit logic, engineers must curate and version the data that defines program behavior, since a model with millions of parameters cannot be patched or reasoned about directly.

Table 1.1: The Paradigm Shift from Software 1.0 to Software 2.0: In Software 2.0, the “programmer” does not write the logic; they curate the dataset that the optimization process uses to write the logic. Debugging therefore moves upstream from code to data. The “compiler” analogy is approximate: unlike a deterministic compiler, the training process is stochastic and may produce different “executables” from the same “source code.”

Feature	Software 1.0 (Traditional)	Software 2.0 (Machine Learning)
Source Code	C++, Python, Java	Training Data + Labels
Compiler	GCC, LLVM	Training loop (stochastic gradient descent)
Logic	Explicit (Hand-coded)	Implicit (Learned)
Failure Mode	Loud (Crash, Exception)	Silent (Metric Degradation)
Debugging	Trace execution path	Inspect data distribution

The data-centered workflow changes what engineers must build and maintain. In a study of production systems at Google, Sculley and colleagues found that model code occupied only a small fraction of the engineering surface (Sculley et al. 2015). The surrounding system—data collection, verification, feature extraction, resource management, serving, and monitoring—was larger and more enduring.

That imbalance creates hidden technical debt. Each surrounding component encodes assumptions about how an example is sampled, what a label means, when a feature is computed, which version reaches serving, and how degradation is detected. None of those assumptions appears in the matrix multiplication that produces a prediction, yet any of them can change the result. Improving the model alone cannot repair a stale proxy, a broken feature pipeline, or a missing feedback signal.

Production data is therefore not a passive input. It is a measurement of the world, collected through a particular product and transformed by a changing pipeline. If that measurement stops representing the intended target, the model can remain numerically healthy while the system becomes wrong. Google Flu Trends provides a revealing example. It had an enormous, timely stream of search data, but the meaning of that data changed beneath the model.



Production ML work is mostly the surrounding system, not the model code alone.

⚠ War Story 1.1: When search logs mistook attention for illness (2014)

Context: Google Flu Trends (GFT) estimated influenza activity from patterns in aggregated search queries (Ginsberg et al. 2009; Lazer et al. 2014). Its attraction came from the speed and scale at which searches arrived relative to clinical reports. If people searched for flu symptoms when they became ill, query volume appeared to offer a timely proxy for infection.

Failure mode: The proxy was not stable. News coverage changed what people searched for, and autocomplete changed how they expressed those searches. Query volume began to measure public attention and product behavior as well as illness. During the 2012–2013 season, GFT estimated roughly twice the CDC-reported proportion of doctor visits for influenza-like illness and overestimated for 100 out of 108 weeks.

Systems lesson: The remedy was not simply a larger search dataset or a better-tuned model. Researchers combined search signals with CDC sentinel clinical data and recalibrated the relationship over time. Data volume is not ground truth: a behavioral proxy needs a feedback loop to a trusted measurement and continual evidence that the proxy still represents the quantity the system claims to estimate.

Google Flu Trends failed without a conventional software defect. The code did not change, but the effective program did because the distribution of inputs and the relationship between proxy and target had changed. Under the **data-as-code principle**, training data does not merely enter a fixed program; it helps determine the operational logic that the model implements. Engineers still write the optimization procedure, but the examples shape the model weights³ through stochastic gradient descent⁴ and related methods. Changing the dataset can therefore change system behavior as surely as changing source code.

From an ML development perspective, this represents a transition from *model-centric* to *data-centric* AI (Ng 2021). In a model-centric approach, teams hold the data fixed and focus on improving model

³ | **Model weights:** The learned numerical parameters of a neural network. A GPT-3-scale (Generative Pre-trained Transformer 3) model stores 175B such values, consuming 350 GB in FP16 precision, a 16-bit floating-point format that uses two bytes per value (Brown et al. 2020). Parameter count determines the weight footprint and strongly influences serving memory traffic and cost (see chapter 5).

⁴ | **Stochastic gradient descent (SGD):** The algorithm learns model parameters by processing small, randomly sampled groups of examples (“batches”) rather than the entire dataset at once. This trades statistical noise for computational speed. Batch size also affects the machine because a batch that is too small may fail to saturate an accelerator’s parallel processors, wasting much of its potential computation.

code. In a data-centric approach, they hold the code comparatively fixed and systematically improve the data, making data curation a first-class part of programming model behavior. The shift also changes what testing can establish because no finite dataset can represent every input that a learned system may encounter.

The limit becomes clear when we ask what a test set can cover. Software 1.0 logic can often be enumerated or partitioned into paths and edge cases. A learned system instead operates across a high-dimensional input space that is technically finite but impossible to enumerate in practice. The 2012 ImageNet challenge makes this problem concrete. AlexNet's decisive win⁵ helped set the deep learning era in motion, yet the benchmark could evaluate only a vanishing fraction of the inputs the model might encounter. Each input is a 224×224 RGB image with $256^{150,528}$ possible pixel configurations, a number with 362,508 digits. The ImageNet Large Scale Visual Recognition Challenge (ILSVRC) validation set contains only 50,000 images (Krizhevsky et al. 2012; Russakovsky et al. 2015). Let **Total Input Space** denote the number of possible inputs and **Test Set Coverage** the number that a test suite actually evaluates. Their disparity creates the **verification gap** (equation 1.1):

$$\text{Verification Gap} = \text{Total Input Space} - \text{Test Set Coverage} \approx \text{Total Input Space} \quad (1.1)$$

The gap does not make testing futile; it changes what testing can establish. Predeployment evaluation provides statistical evidence over sampled inputs, while production monitoring tests whether the operating population and observed outcomes still support that evidence. Statistical reliability replaces any expectation of exhaustive proof.

The engineering consequence reaches beyond testing. Debugging an ML system requires debugging the data, not the Python scripts alone. Version control must track datasets, not git commits alone. Evaluation must examine distributions and outcomes, not code paths alone. Together, these practices make learned behavior traceable, but they cannot turn statistical evidence into exhaustive proof.

Input space

Test set

Test coverage is a vanishing fraction of the input space.



Checkpoint 1.1: The paradigm shift

Before tracing the history of AI, verify your understanding of the paradigm shift in how we build software:

- ☐ **Software 1.0 vs. 2.0:** can you distinguish Software 1.0 (explicit instructions) from Software 2.0 (optimization objectives)?
- ☐ **Debugging shift:** can you identify which dataset evidence to inspect when accuracy falls without a code change?
- ☐ **Verification gap:** can you explain why finite test coverage cannot establish statistical reliability across the full production input distribution?

The verification gap marks a deeper change in what engineers can claim. Traditional assertions often describe a particular execution in which a given input and state lead the program down a specific path. ML performance claims describe behavior over a population. A fixed model may return the same output for the same input while its measured accuracy changes as the population changes. Data supplies the patterns from which the system learns, but its noise, drift, and omissions also create uncertainty. Robustness therefore cannot mean resisting every change. It requires making change observable and adapting when the evidence no longer supports the system's assumptions, a form of **probabilistic engineering**.

Learning behavior from data can now seem inevitable, but it was not. Earlier AI systems attempted to program intelligence directly. Following those attempts reveals why the bottleneck moved from logic, to knowledge, to features, and finally to infrastructure—and why systems engineering became central to progress.

1.3 The Evolution of AI Bottlenecks

AI's evolution reveals a progression of bottlenecks, each overcome by systems innovations that expanded what was computationally possible. One foundational landmark is Turing's⁶ paper

⁵ | **AlexNet:** Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton trained this convolutional neural network across two GPUs. At ILSVRC 2012, its 15.3 percent top-5 error, meaning the correct class was absent from its five highest-scoring predictions, substantially beat the second-place system's 26.2 percent (Krizhevsky et al. 2012). Section 1.3.3 returns to its architecture and systems co-design.

⁶ | **Alan Turing:** His 1950 "Imitation Game" reframed intelligence as an output-measurement problem: judge a system by what it *does*, not by what it *is*. This engineering-first stance persists in every ML systems metric we use today: accuracy, latency, throughput, and FLOP/s per watt are all output measurements. The iron law (section 1.7) decomposes performance into observable, measurable terms rather than internal architectural properties for exactly this reason.




Figure 1.1: AI Development Timeline: A chronological proxy curve traces historical mentions of artificial intelligence in published literature (Michel et al. 2011), with light-brown bands marking the two AI Winter periods (1974–1980, 1987–1993). Callout boxes mark Turing’s 1950 paper (Turing 1950), the Dartmouth conference (McCarthy et al. 1955), the Perceptron (Rosenblatt 1958), ELIZA (Weizenbaum 1966), Deep Blue (Campbell et al. 2002), and GPT-3 (Brown et al. 2020).

“Computing Machinery and Intelligence” (Turing 1950), which asked whether machines can think. Early systems explored contrasting approaches: the Perceptron (1958) (Rosenblatt 1958) learned from examples, while ELIZA⁷ (Weizenbaum 1966) followed hand-written pattern-matching rules. Both remained narrow, but for different reasons. Subsequent eras hit the knowledge acquisition bottleneck because manual knowledge entry could not scale. Modern systems face a different constraint in computational throughput.

The timeline in figure 1.1 traces how often artificial intelligence is mentioned in published books, a proxy for attention rather than a direct measure of research output. It reveals a recurring pattern of intense optimism followed by “AI winters”⁸ when funding collapsed, often after systems limitations exposed a gap between ambition and available capability. Resurgences combined algorithmic advances with new data and engineering infrastructure. Each resurgence displaced one bottleneck and exposed the next.

1.3.1 The prelearning era: Logic and knowledge bottlenecks

Before machine learning existed as a discipline, engineers attempted to build intelligent systems through two successive paradigms, each of which hit a fundamental scaling barrier. Symbolic AI encoded intelligence as logical rules and hit the *logic bottleneck* when those rules could not capture real-world ambiguity. Expert systems encoded intelligence as domain knowledge and hit the *knowledge bottleneck* when acquiring and maintaining that knowledge became more expensive than the systems were worth. Together, these two eras reveal the pattern that motivates everything that follows. Hand-crafted representations do not scale.

⁷ | **ELIZA:** A 1966 natural-language program using pattern-matching rules rather than learned parameters. Its scripts could store and retrieve selected inputs, but this limited hand-written memory did not provide learned conversational state. Every new input variation required another rule, making maintenance grow faster than capability and foreshadowing the knowledge bottleneck that constrained later expert systems.

8 | AI winters as systems

failures: The first AI winter (1974–1980) unfolded amid funding cuts; the 1973 Lighthill Report criticized the gap between AI promises and delivered results (Lighthill 1973). The second winter (1987–1993) involved a market and funding collapse around expert systems and specialized Lisp machines as general-purpose workstations undercut their economics (Hendler 2008). From this book's systems perspective, both episodes expose algorithm ambition outrunning available infrastructure, market support, and engineering maturity, not merely a shortage of clever algorithms.

9 | Dartmouth Conference

(1956): The workshop organized around the term “artificial intelligence,” already used in its 1955 proposal (McCarthy et al. 1955). Its participants framed intelligence in terms of language, abstraction, problem solving, and self-improvement, with little attention to the physical constraints of storage and compute that later became central. The same compute-agnostic assumption, that a better algorithm could always overcome a hardware limit, is precisely what this book exists to correct: every chapter that follows argues that systems constraints are first-class design variables, not afterthoughts.

10 | STUDENT: Daniel Bobrow's 1964 MIT program parsed constrained English word problems, represented their relationships symbolically, and passed the resulting equations to an algebra solver. It was an early separation of language interpretation from formal reasoning: once the representation was correct, solving was straightforward; coverage depended on hand-written transformations (Bobrow 1964).

1.3.1.1 The symbolic AI era and the logic bottleneck

The first era of AI engineering (1950s–1970s) attempted to reduce intelligence to **symbolic AI** manipulation, an approach later crystallized in the physical-symbol-system hypothesis (Newell and Simon 1976). Researchers at the 1956 Dartmouth Conference⁹ (McCarthy et al. 1955) hypothesized that aspects of intelligence could be precisely described and simulated by machines. Even then, Arthur Samuel at IBM demonstrated a different path in 1959 when a checkers program improved through self-play. His work coined the term “**machine learning**” (Samuel 1959), though the dominant paradigm remained symbolic. Daniel Bobrow's STUDENT¹⁰ system exemplifies this approach (Bobrow 1964).

These systems could produce impressive demonstrations, yet they were operationally brittle. Their success depended on manually coded rules that mapped each acceptable input form to a symbolic representation. STUDENT could solve an algebra problem once it translated English into equations, but a minor variation in phrasing could prevent that translation. The difficulty extended beyond language. Hans Moravec's¹¹ work on autonomous navigation at Stanford revealed that tasks humans find trivial (seeing, walking, grasping) were far harder to engineer than tasks humans find difficult, like chess or algebra.

Example 1.1: STUDENT (1964)

STUDENT separated language interpretation from algebra:

```
language "Two numbers sum to 30; one is twice the other."
parse    x + y = 30, x = 2y
solve    x = 20, y = 10
```

Once the representation was correct, solving was routine. The bottleneck was the translation: each unfamiliar phrasing needed another hand-written rule, so linguistic coverage grew only as fast as the rule base.

1.3.1.2 The expert systems era and the knowledge bottleneck

In the expert-systems era, engineers narrowed the problem. Instead of constructing general reasoning systems, they encoded deep knowledge from a specific domain. MYCIN, designed to diagnose blood infections, let medical knowledge be expressed as production rules (Shortliffe et al. 1975).

Example 1.2: MYCIN (1976)

MYCIN made its reasoning explicit. Patient facts triggered readable production rules:

```
facts  stain=gram-positive, shape=coccus, arrangement=clumps
rule   IF facts match THEN organism=staphylococcus (certainty=0.7)
cycle  match facts -> fire rule -> update certainty -> repeat
```

An inference engine repeated this cycle across hundreds of rules and propagated certainty factors toward a diagnosis. A physician could inspect every step, but coverage still grew one rule at a time: each exception had to be elicited, encoded, and reconciled with the existing rule base.

MYCIN performed well in specific tests, but its success exposed the **knowledge acquisition bottleneck**.¹² The problem was no longer whether a rule could express expertise. It was how tacit judgment entered the system and remained consistent as the rule base grew.

Faster machines could evaluate more rules, but they could not extract human judgment faster. Scalable AI needed a way to infer useful decision boundaries from examples rather than requiring engineers to enumerate those boundaries in advance.

That change did not remove human design; it changed what engineers designed. Instead of encoding every decision, they chose the examples, representations, objectives, and measurements

from which a decision could be learned. The next era therefore exchanged the knowledge-acquisition bottleneck for a new question: which evidence should the learner see?

1.3.2 The statistical learning era and the feature engineering bottleneck

The 1990s marked the shift to **statistical learning** and probabilistic systems. Instead of hard-coded logic, systems estimated probabilities from data ($p(y | x)$). This transition was driven by the availability of digital data and the “unreasonable effectiveness”¹³ of large datasets.

Spam filtering illustrates this shift. Rather than maintaining lists of forbidden words, statistical filters learned the probability that a word implies spam based on millions of examples.

Example 1.3: Early spam detection systems

Statistical filters shifted the decision from a hand-written rule to a probability estimated from labeled messages:

```
rules: if "free" or "winner" appears -> spam
learn: labeled email -> word likelihoods -> P(spam | email)
```

Maintenance did not disappear. It moved from editing keyword lists to curating representative examples and checking whether the learned probabilities still tracked current email.

Learning the decision boundary from data removed one bottleneck but exposed another. Statistical algorithms such as Support Vector Machines (SVMs) could learn robustly only after humans converted raw inputs into structured features. The learner adjusted the boundary; engineers decided which evidence it could see. Scaling to a new problem often meant rebuilding the preprocessing stack, turning an apparent algorithm limitation into the **feature engineering bottleneck**. The traditional pipeline makes this manual effort visible because several hand-crafted stages preceded any learning at all.

This hybrid approach combined human-engineered features with statistical learning. The Viola-Jones algorithm¹⁴ (Viola and Jones 2001) exemplifies this era, achieving real-time frontal-face detection using simple rectangular features and cascaded classifiers. It showed that well-engineered features could enable practical low-latency applications, but only within narrow domains where experts could hand-craft the right representations.

Example 1.4: Traditional computer vision pipeline

A conventional vision system separated representation from classification:

```
image -> resize -> HOG/SIFT features -> SVM -> label
```

Only the final decision boundary was learned. Moving from faces to pedestrians or another domain meant redesigning the feature extractor and often the entire preprocessing stack.

1.3.3 The deep learning era and the infrastructure bottleneck

Deep learning changed which part of the system learned. Instead of receiving features designed by humans, neural networks learned representations directly from raw inputs such as pixels and audio waveforms. The training process could now shape both the representation and the decision boundary, enabling “end-to-end” learning.

The breakthrough was not algorithmic alone. Convolutional neural networks (CNNs) existed earlier (LeCun et al. 1998, 2015); AlexNet paired architecture and training with **systems co-design**, choosing the model, training procedure, and hardware mapping together (Krizhevsky et al. 2012). Its parallel matrix operations matched GPU capabilities. With 60 million parameters distributed across two GTX 580 GPUs, AlexNet achieved 15.3 percent top-5 error, a 41.6 percent relative improvement over the next-best entry that year. The processing stages in figure 1.2 show the model learning progressively richer image representations before producing one of 1,000 output classes.

¹¹ | **Moravec’s paradox:** Carnegie Mellon roboticist Hans Moravec observed that high-level reasoning (chess) requires little compute while low-level perception (walking) requires massive parallelism (Moravec 1988). This paradox explains a central fact of ML systems engineering: the tasks that seem “easy” to humans (vision, speech, motor control) are the ones that demand the highest FLOP/s, memory bandwidth, and specialized hardware, driving the accelerator revolution that defines modern ML infrastructure.

¹² | **Knowledge acquisition bottleneck:** Feigenbaum’s knowledge-engineering work framed applied AI around the practical difficulty of extracting, representing, and maintaining expert knowledge (Feigenbaum 1984). In systems terms, this bottleneck was a throughput problem: knowledge elicitation and rule maintenance were bound by the serial bandwidth of human experts. Unlike computational bottlenecks that yield to faster hardware, this one was the original “does not scale” constraint in AI and a direct motivation for the data-driven paradigm that followed.

¹³ | **Unreasonable effectiveness of data:** The observation that a simple statistical model fed with massive amounts of data can outperform a more sophisticated model with less data (Halevy et al. 2009). Halevy and colleagues illustrated the effect with large web corpora, not a universal error-rate multiplier; gains depend on the task, model, data quality, and starting point. The result supported the shift from brittle hand-crafted systems toward probabilistic models and made data collection, storage, pre-processing, and distributed training central engineering concerns.




Figure 1.2: AlexNet Architecture: The network that launched the deep learning revolution at ImageNet 2012 (Krizhevsky et al. 2012). Two parallel GPU streams process 224×224 input images through convolutional layers (green blocks) that extract spatial features at decreasing resolutions, converging through three fully connected layers to 1,000 output classes. With 60 million parameters trained across two GTX 580 GPUs, AlexNet achieved 15.3 percent top-5 error, a 41.6 percent relative improvement over the second-place entry.

14 | Viola-Jones algorithm: The algorithm’s real-time speed came from a classifier cascade that used simple, hand-engineered rectangular features to immediately reject nonface regions. The method was designed and evaluated for frontal-face detection, illustrating the era’s trade-off: expert feature design could be fast and effective, but the representation was task-specific. The first two layers alone could discard over 80 percent of negative sub-windows while using just twelve of the 6,000+ total features (Viola and Jones 2001).

Hardware still shaped what could be learned. With only 3 GB of Video Random-Access Memory (VRAM) per GTX 580, AlexNet divided convolutional and dense layers between two GPU streams. This early model-parallel design foreshadowed modern multi-GPU training. Deep learning reduced the need for hand-crafted features, but it moved the binding constraint to the infrastructure required to store data, move parameters, and coordinate computation.

Deep learning effectively traded the feature engineering bottleneck for a new **compute bottleneck**. Models like GPT-3 (175 billion parameters) illustrate the scale of this new challenge. Brown et al. (2020) report training on about 300 billion tokens from filtered web text, books, and Wikipedia. Using the book’s dense-training approximation, that parameter-token scale implies roughly 314 zettaFLOPs of compute (1 zettaFLOP = 10^{21} FLOPs). The token dataset itself occupies roughly 420 GB. Because the original paper does not specify the exact hardware cluster, translating this into accelerator-years serves as an illustrative systems estimate rather than a recorded benchmark. The primary engineering challenge shifted from describing a cat’s ear to coordinating large-scale distributed training without failure.

Each major transition in AI history traces back to a shifting physical constraint rather than algorithmic invention alone. Table 1.2 compares the four major eras across their core reasoning strengths, binding systems bottlenecks, and operational data requirements.

Table 1.2: AI Paradigm Evolution: Each era is defined by the systems bottleneck that constrained it. Deep learning (far right) overcame the Feature Engineering bottleneck but introduced new infrastructure challenges, necessitating modern ML systems engineering.

Aspect	Symbolic AI	Expert Systems	Statistical Learning	Deep Learning
Key Strength	Logical reasoning	Domain expertise	Versatility	Pattern recognition
Bottleneck	Brittleness (Rules break)	Knowledge Entry (Experts are scarce)	Feature Engineering (Manual preprocessing)	Compute & Data Scale (Infrastructure cost)
Data Handling	Minimal data needed	Domain knowledge-based	Moderate data required	Massive data processing

15 | Richard Sutton: A reinforcement learning pioneer whose 2019 essay crystallized the pattern traced in section 1.3: from symbolic AI through expert systems to deep learning, general methods using computation consistently outperformed hand-engineered expertise. The lesson is “bitter” because it implies that domain-specific logic is a depreciating asset, while the durable advantage belongs to systems engineering that can absorb the billion-fold increase in raw compute since the 1970s.

1.4 The Bitter Lesson

Expert systems invested engineering effort in encoding domain knowledge; deep learning systems invest that effort in absorbing more data and computation. The **bitter lesson** captures the historical pattern in which general methods that use increasing computation consistently outperform approaches that encode human expertise. Richard Sutton¹⁵ crystallized this insight in his 2019 essay “The Bitter Lesson” (Sutton 2019). Sutton wrote, “The biggest lesson that can be read from 70 years of AI research is that general methods that leverage computation are ultimately the most effective, and by a large margin.”

Table 1.3 pairs representative benchmark milestones with their hardware substrates. Its final column traces the progression from single-threaded CPU rule evaluation to multi-thousand-accelerator clusters, using 2.5 million reference GPU-days as an illustrative frontier-training anchor (Patel and Wong 2023).

Table 1.3: AI Performance Evolution Across Paradigms: Representative systems illustrate the growth in computational scale across eras. GPT-4 training details are estimated reference anchors rather than official disclosures.

Era	Approach	Representative Task	Performance	Computational Resources
Expert Systems (1980s)	Hand-crafted rules	Chess (Elo rating)	System-dependent	Minimal (rule evaluation)
Statistical ML (1990s–2000s)	Feature engineering + learning	Handwritten digit recognition	about 98–99% on MNIST-era benchmarks (LeCun et al. 1998)	CPU-era feature pipelines; resources varied by implementation
Deep Learning (2012)	End-to-end neural networks	ImageNet top-5 accuracy	84.7% (AlexNet)	6 days on 2 GPUs
Modern Deep Learning (2020+)	Large-scale transformers	ImageNet top-1 accuracy	88.55% (ViT-H/14) (Dosovitskiy et al. 2021)	Large-scale Tensor Processing Unit (TPU) pretraining
Modern Deep Learning (2023)	Foundation models	MMLU benchmark	86.4% (GPT-4) (OpenAI et al. 2023)	Estimated ~2.5 million reference GPU-days (on the order of 25,000 reference GPUs run for 90 days) (Patel and Wong 2023)

ImageNet top-5 accuracy and Massive Multitask Language Understanding (MMLU) (chapter 12) measure distinct capabilities, but share a systems trajectory. Hardware scale expanded from single-node CPU execution to multi-megawatt clusters, corroborating Sutton’s observation that methods designed to use raw computation consistently outpace hand-tuned representations.

The principle finds further validation across AI breakthroughs. In chess, IBM’s Deep Blue defeated world champion Garry Kasparov¹⁶ in 1997 by combining custom chess hardware, large-scale search, and chess-specific evaluation knowledge. Its evaluation function encoded human chess heuristics, but the scale of search enabled by custom silicon was central to turning that knowledge into championship-level play. In Go, DeepMind’s AlphaGo¹⁷ (Silver et al. 2016) achieved superhuman performance by combining supervised learning from expert games with reinforcement learning through self-play and neural-network-guided tree search, rather than relying on hand-coded Go strategy.

The lesson is “bitter” because our intuition misleads us. We naturally assume that encoding human expertise should be the path to artificial intelligence. Yet repeatedly, systems that use computation to learn from data outperform systems that rely on human knowledge given sufficient scale. The pattern has held across symbolic AI, statistical learning, and deep learning eras.

Modern language models like GPT-4 and image generation systems like DALL-E illustrate this principle directly. Their capabilities emerge not from linguistic or artistic theories encoded by humans but from training general-purpose neural networks on vast amounts of data using substantial computational resources. Estimates for models at GPT-3’s scale suggest roughly 1.3 GWh of energy¹⁸ (Patterson et al. 2021), and serving these models to millions of users turns inference into a continuous data-center power, cooling, and capacity-planning problem.

The implication is that realizing the bitter lesson’s promise requires expertise in data engineering, hardware optimization, and systems coordination¹⁹ that goes far beyond algorithmic innovation. Chapter 11 quantifies these constraints once the necessary foundations are in place, including the memory bandwidth limits that shape system design.

The world is rushing to build AI systems, but building them is not enough. Sutton’s bitter lesson shows why. Progress depends on the data, computation, and infrastructure that make learning at scale possible. This book therefore treats the machine learning system as an engineering object: one that must be designed, deployed, evaluated, and maintained under real constraints. That broader responsibility is AI engineering, the discipline defined formally in section 1.9. Before defining the discipline, we must first define its object, beginning with a familiar production workload.

1.5 Defining ML Systems

Return to the spam filter introduced in the history of statistical learning. At production scale, the same apparently simple classifier operates against global email traffic measured in hundreds of billions of sent and received messages per day (Statista Research Department 2024), and large

¹⁶ | **Deep Blue:** IBM’s chess system (Campbell et al. 2002) defeated World Champion Garry Kasparov in 1997 through a systems combination: search at roughly 200 million positions per second on 480 custom chess processors, plus chess-specific evaluation and knowledge. Deep Blue was an early public demonstration that purpose-built silicon could amplify search and encoded domain knowledge, foreshadowing the domain-specific accelerator strategy that defines modern ML hardware.

¹⁷ | **AlphaGo:** AlphaGo first learned from human expert games, then improved through reinforcement learning from self-play, trading hand-coded Go strategy for a data-and-compute pipeline that could explore the problem space at massive computational scale. After three days of self-play training, AlphaGo Zero surpassed the original AlphaGo, winning 100 games to 0 (Silver et al. 2017).

¹⁸ | **GPT-3 (Generative Pre-trained Transformer 3) training energy:** Patterson et al. (2021) estimated GPT-3’s single training run consumed approximately 1,287 MWh and emitted 552 tonnes of CO₂-equivalent, roughly the annual electricity of 120 average US households using a 10.7 MWh/household-year baseline. The energy cost is shaped not only by arithmetic but also by data movement through the memory hierarchy; moving data across memory levels can cost orders of magnitude more energy than local arithmetic (Horowitz 2014).

19 | Memory bandwidth:

The rate at which a model's parameters move from memory to the processor. The gigawatt-hour-scale energy consumed by GPT-scale training is shaped not only by computation but also by the physically expensive process of fetching billions of weights through the memory hierarchy. Moving data from off-chip memory can cost one to several orders of magnitude more energy than local arithmetic, depending on precision and memory level, making bandwidth, not processor speed alone, a direct driver of the data center's massive power draw.

providers must decide in milliseconds which messages deserve attention and which should be quarantined.

This deceptively simple task reveals *what* distinguishes machine learning systems from traditional software. The challenge begins with data. The filter trains on millions of labeled examples and must keep adapting as spammers evolve their tactics, rather than relying on programmers to encode every spam pattern manually. It then becomes an algorithmic problem, because the model must generalize from those examples to messages it has never seen before while balancing precision against recall so legitimate email is not hidden. Finally, the same decision becomes an infrastructure problem. Providers must process billions of emails daily, store and update models as spam evolves, and serve predictions with sub-100 ms latency across horizontally scaled data centers. The classifier is therefore only one component of a changing data, software, and infrastructure system. That observation gives us the object this book studies.

When a new phishing template appears, the data layer must capture representative messages, the algorithm must separate attacks from legitimate mail, and the machine must distribute updated parameters before delivery. Any layer can fail while the others continue working, so the full chain from measurement through learning to execution is the right unit of analysis.



Definition 1.1: Machine learning systems

Machine learning systems are software systems whose core behavior is determined by parameters learned from data rather than explicitly programmed rules, making performance a function of data quality, algorithm choice, and hardware capacity simultaneously.

1. **Significance:** Requirements are end to end. A spam filter is useful only when its training data represents current attacks, its model separates malicious from legitimate mail, and its serving path classifies each message within the delivery budget. Failure in any layer changes the user-visible result.
2. **Distinction:** Unlike traditional software, an ML system's accuracy can change when the world changes even if its code and weights do not. The production input distribution may move relative to what the model learned, silently changing performance without an error or exception.
3. **Common pitfall:** The model is not the system. Data pipelines, feature transformations, serving infrastructure, monitoring, and feedback loops surround the learned parameters and often dominate the engineering burden (Sculley et al. 2015).

The definition exposes three diagnostic axes: the data that describes the task, the algorithm that turns those examples into behavior, and the machine that executes that behavior within the operating budget. The D·A·M taxonomy gives those axes a reusable form.

The same user-visible symptom can originate in any of the three. A missed phishing message may mean that representative examples were absent, that the learned decision boundary was inadequate, or that the serving path missed its latency budget. Treating every miss as a model problem risks improving the wrong component. A useful systems taxonomy must separate these causes without pretending that they are independent.

Each explanation calls for different evidence. The data hypothesis sends the engineer to coverage, labels, and distribution shift. The algorithm hypothesis sends the engineer to error slices, model capacity, and the learning objective. The machine hypothesis sends the engineer to latency, throughput, memory, and utilization. These investigations are not interchangeable. Faster hardware cannot supply examples that were never collected, while a larger dataset cannot repair a serving path that misses its deadline because memory is saturated.

The relevant constraint is the one whose relaxation improves the end-to-end result. This idea of a *binding constraint* prevents teams from optimizing the most visible component instead of the component that governs the outcome. The purpose of D·A·M is therefore operational: it identifies the next hypothesis to test and the class of intervention capable of changing system behavior. The diagnosis is provisional. Once an intervention relaxes one limit, the system must be measured again because a different axis may now bind. D·A·M is a loop, not a one-time label.

At first pass, D·A·M is a three-part mental model. Data supplies evidence, Algorithm turns it into behavior, and Machine executes it within the operating budget. The axes describe the system as a whole, not isolated pipeline stages.

Definition 1.2: The D·A·M taxonomy

D·A·M taxonomy is a diagnostic framework that classifies any machine learning system performance bottleneck along three axes. Data determines what examples and bytes the system must process, Algorithm determines the model structure and work required to learn or predict, and Machine determines the hardware capacity available to execute that work. The goal is to identify which axis is the binding constraint.

1. **Significance:** The diagnostic power is concrete even before detailed hardware arithmetic enters the story. If the spam filter misses a new phishing campaign because the training set never contained that tactic, the binding axis is Data. If the training examples are adequate but the model cannot express the pattern, the binding axis is Algorithm. If both are adequate but the service cannot classify messages quickly enough during a traffic spike, the binding axis is Machine. Quantitative diagnosis begins by asking which axis is limiting the system.
2. **Distinction:** Unlike traditional software performance analysis, which treats code and data as separate concerns, the D·A·M taxonomy recognizes that algorithm choice directly determines both the training dataset size required (a transformer needs orders of magnitude more data than a linear model to generalize) and the machine required to run it.
3. **Common pitfall:** A frequent misconception is that the three axes are independent. Changing from a simple classifier to a larger model can require more memory, different serving infrastructure, and a broader data distribution. The axes move together.

Figure 1.3 turns the definition into a first-pass map.




Figure 1.3: The D·A·M Axis Triad: Data, Algorithm, and Machine are three interdependent parts of an ML system. The bidirectional links emphasize that changing one part can change the requirements imposed on the other two.

The triangle provides the first diagnostic pass. It converts a visible outcome into three measurement paths: inspect the evidence, characterize the work, and measure the execution. Each path demands different evidence. Coverage and distribution measurements probe Data, error patterns and operation counts probe Algorithm, and latency, bandwidth, and utilization probe Machine. The diagnostic process begins with an end-to-end result, such as accuracy, latency, throughput, or cost, and tests each path rather than assuming that the component nearest the symptom is responsible. A diagnosis earns confidence only when relaxing its candidate constraint improves the full-system result. The arrows matter because an intervention at one vertex can move the bottleneck to another.

The triangle deliberately suppresses those interactions to keep the first diagnosis simple. Once a candidate axis is identified, the next step is to trace what that constraint depends on and what it constrains. Real systems rarely remain at one vertex. Data formats change memory traffic, model structure changes hardware utilization, and deployment constraints reshape both data and algorithm choices. Figure 1.4 expands the triangle into the richer design space created by these intersections.




Figure 1.4: The D·A·M Intersection Landscape: Data, Algorithm, and Machine begin as distinct diagnostic lenses, but practical ML systems live at their boundaries. Pairwise intersections ask what a system can learn from, how information moves, and how computation executes efficiently. Their center is ML systems engineering, where all three constraints must be balanced together.

Read the Venn diagram from the outer regions inward. Data and Algorithm determine *what the system can learn from*. Data and Machine determine *how information moves*. Algorithm and Machine determine *how computation executes efficiently*. At the center, all three questions must be answered together. That center is the discipline of ML systems engineering.

The overlap labels are also signposts for the chapters ahead. Several names may be unfamiliar because they denote problems the rest of the book develops in sequence. Data selection and model training develop the Data–Algorithm question. Data engineering and hardware acceleration trace the Data–Machine path. Frameworks, compression, and serving develop the Algorithm–Machine question. The figure is therefore an orientation, not a list to memorize. Later chapters turn each region into concrete measurements, design choices, and interventions.

The techniques named in the figure also show why no choice remains local: data selection changes the work presented to the algorithm, architecture changes memory and compute demand, and machine choice changes which algorithms and data paths are practical. Throughout this book, first locate the binding region, then optimize.

The D·A·M landscape provides the diagnostic lens, but building systems also requires a layered view that connects physical limits to a user-facing mission.

1.5.1 The four-layer hierarchy from silicon to mission

Every machine learning system analyzed in this text is constructed from four hierarchical layers, ensuring that a decision made at the silicon level is traceable to its impact on the final mission.

1. **Hardware (The Silicon).** The physical foundation (The Engine) defines peak compute throughput (R_{peak}), memory bandwidth (BW), and memory capacity. Concrete hardware twins instantiate those quantities when deployment scenarios need numeric constraints.
2. **Systems (The Platforms).** The integrated deployment unit (The Car) defines the operating envelope through its power budget, thermal limits, and node-level interconnects. Examples include the Training Cluster Node or the Sub-Watt Sensor Node.

3. **Workloads (The Models).** The algorithmic demand (The Route) comprises the operation count (O), data volume moved (D_{vol}), and data layout. Scenario-specific workloads, such as GPT-4 and Wake Vision (a visual wake-word dataset sized for microcontrollers), instantiate these demands for particular missions.
4. **Missions (The Scenarios).** The application context (The Destination) occupies the top of the stack, where a system is deployed to solve a specific problem. A mission introduces requirements such as battery life, safety latency, or cloud cost ceilings that dictate the configuration of every layer below.

This hierarchy ensures that when we build a lab or a case study, engineers are not starting from scratch, but rather inheriting the constraints of a deployment paradigm and applying a scenario workload to a specific mission. The lifecycle discussion in section 1.9.1 pairs each recurring mission with its workload and binding constraint. This structured approach allows us to reason about the “Physics of ML” across any application domain.

Systems Perspective 1.1: The ML systems landscape: Four deployment paradigms

The machine learning systems landscape spans roughly 10^6 in computational power and 10^5 in memory capacity. Table 1.4 contrasts the memory, compute, and power envelopes across the four **deployment paradigms** that define the constraints for every subsequent chapter.

Table 1.4: Four Deployment Paradigms: Representative memory, compute, and power envelopes for cloud, edge, mobile, and TinyML systems, exposing the multi-order-of-magnitude span that prevents simple model reuse across tiers. Exact hardware twins and peak-rate calculations appear in the hardware chapters.

Paradigm	Representative System	Memory Envelope	Compute Envelope	Power Envelope
Cloud	Data-center accelerator node	Large device memory plus storage ($\approx 10^{11}$ B)	Highest-throughput tier ($\approx 10^{15}$ ops/s)	Facility-managed power
Edge	Robotics or industrial gateway	Local memory under deployment limits ($\approx 10^{11}$ B)	Local accelerator or CPU budget ($\approx 10^{14}$ ops/s)	Wall, vehicle, or site power
Mobile	Smartphone or wearable-class SoC	Shared application memory ($\approx 10^{10}$ B)	Phone-class neural, GPU, and CPU engines ($\approx 10^{13}$ ops/s)	Battery and thermal cap
TinyML	Microcontroller node	Kilobyte-scale memory ($\approx 10^6$ B)	Always-on sensor compute ($\approx 10^9$ ops/s)	Milliwatt-class battery budget

Reading the memory and compute columns from Cloud to TinyML shows endpoints that differ by 10^5 in memory and 10^6 in compute. This divergence is precisely why engineers cannot shrink a cloud model to run at the edge; each tier requires a fundamental redesign of the D·A·M axes.

The D·A·M taxonomy serves as a diagnostic lens throughout this text. Scale in ML systems is the relentless pursuit of the moving bottleneck. Alleviating a constraint along one axis often shifts the limitation to another. Upgrading to faster GPUs (Machine) might reveal that storage cannot feed data fast enough (Data). Collecting a massive dataset (Data) might reveal that the model lacks capacity to learn from it (Algorithm). Switching to a larger model (Algorithm) might exceed available memory (Machine). The recurring engineering task is to identify the axis that binds, intervene there, and check whether the intervention has merely created a worse bottleneck elsewhere.

The multi-order-of-magnitude span across these four paradigms is not merely a technical curiosity; it translates directly into cost. A model that fits comfortably in a data-center accelerator’s memory cannot run unchanged on a microcontroller-class device, and bridging that gap requires engineering trade-offs at every tier of the D·A·M taxonomy. Data quality, algorithmic efficiency, and hardware capability interact through the economic constraint of **useful work per dollar**.

🔍 Systems Perspective 1.2: Useful work per dollar

Researchers often optimize for *accuracy*, whereas systems engineers must optimize for *useful work per dollar*. Let O_{total} denote the total operations performed across all sample presentations, and let cost efficiency denote useful operations delivered per dollar. Equation 1.2 provides a dimensionally consistent first-order cost model:

$$\text{Compute Cost} \approx \frac{O_{\text{total}}}{\text{Useful Operations per Dollar}} \quad (1.2)$$

Each D·A·M axis improves this ratio differently. Better data can reduce the sample presentations needed to reach a target quality. More efficient algorithms reduce the operations required per sample, while more cost-efficient hardware increases the useful operations delivered per dollar.

Systems engineering is the art of balancing this equation. A 10 percent gain in cost efficiency can fund about 10 percent more useful operations at the same budget, but whether to spend them on more data, more training passes, or a larger model depends on the workload's learning-curve elasticity. If error scales approximately as $D^{-\alpha}$ for dataset size D , the gain from more data is governed by $\alpha \log(1.1)$ rather than by a universal percentage. The engineer's job is to estimate that elasticity for the system at hand and decide whether the trade-off is economically viable.

These interactions change more than cost; they also change how failure appears. A data shortcut, model change, or hardware bottleneck can all surface as degraded behavior after deployment, often without a crash or exception.

1.6 ML vs. Traditional Software

The D·A·M taxonomy reveals that ML systems comprise data that guides behavior, algorithms that extract patterns, and machines that enable learning and inference.²⁰ To see why engineering them differs from traditional software, follow that failure into production.

Many software defects produce explicit failure modes. Applications crash, error messages propagate, and monitoring systems trigger alerts, enabling rapid diagnosis and remediation. Conventional software can also return incorrect results silently, but machine learning adds **silent degradation**, in which performance declines without triggering conventional error detection mechanisms. The algorithms continue executing and the machines maintain prediction serving, yet the learned behavior becomes progressively less accurate or contextually relevant.

Return to the driver-assistance systems from the opening. Conventional automotive software often exposes faults through diagnostic warnings, although it can also produce silent errors. An ML-based perception system adds a different challenge. Its accuracy in detecting pedestrians might decline from 95 percent to 85 percent over several months as seasonal changes introduce lighting conditions, clothing patterns, or weather phenomena underrepresented in training data. The vehicle continues operating, successfully detecting most pedestrians, yet the degraded performance creates safety risks that become apparent only through systematic monitoring of edge cases and comprehensive evaluation. Conventional error logging and alerting mechanisms remain silent while the system becomes measurably less safe.

The magnitude of this degradation matters in safety-critical contexts. A perception model running at 10 Hz processes 36,000 frames in one hour. A 0.1 percent false-negative rate applies only to relevant positive cases; the miss count also depends on their frequency and on temporal filtering, sensor fusion, and operational-design-domain limits. The 10-percentage-point degradation from 95 percent to 85 percent is therefore not merely an accuracy change; it changes the exposure rate of downstream control logic in precisely the edge cases where detection was already marginal.

This silent degradation manifests across all three D·A·M axes. The data distribution shifts as user behavior evolves, seasonal patterns emerge, and new edge cases appear (Gama et al. 2014; Quiñonero-Candela et al. 2009). Meanwhile, the algorithms continue making predictions based on outdated learned patterns, unaware that their training distribution no longer matches operational reality. The machines faithfully serve these increasingly inaccurate predictions at scale, amplifying the problem across every user and every query.

²⁰ | **Inference:** From Latin *inferre* ("to bring in" or "to conclude"). In ML engineering, inference refers to the deployment phase where a trained model applies learned patterns to novel inputs. The systems distinction matters because training is throughput-optimized (maximize samples/second), whereas inference is latency-optimized (minimize milliseconds/prediction). These opposing objectives demand fundamentally different hardware configurations and software stacks (see chapter 13).

Because this failure mode is silent, traditional crash logs cannot detect it; quantitative signals are needed to test whether a measured distribution shift predicts performance loss. Just as hardware execution time can be decomposed into physical constituents, reliability degradation can be modeled as a function of environmental change. Here, Accuracy_0 is initial accuracy at deployment, $\mathcal{D}(P_t \| P_0)$ is the statistical divergence between current distribution P_t and training distribution P_0 , and λ is locally fitted sensitivity to the chosen shift measure. This relationship, stated in equation 1.3, is the **degradation equation**. It is a local diagnostic approximation that illustrates when greater distribution shift accompanies performance loss over time, not a universal prediction law.

$$\text{Accuracy}(t) \approx \text{Accuracy}_0 - \lambda \cdot \mathcal{D}(P_t \| P_0) \quad (1.3)$$

This first-order linearization can capture a local trend when labeled observations support the relationship. **Data drift** occurs when P_t differs from P_0 , so predictions can become unreliable even when the code has not changed. The model can break down for large shifts, and the divergence measure $\mathcal{D}(\cdot \| \cdot)$ remains deliberately general (common choices include KL divergence, total variation distance, or Wasserstein distance). Divergence alone does not determine the sign or magnitude of an accuracy change. With these limitations, the diagnostic suggests three engineering levers.

1. **Improve initial accuracy** (Accuracy_0). Better training, more data, and superior architectures shift the curve but not its slope.
2. **Reduce distribution sensitivity** (λ). Robust training techniques, domain adaptation, and broader training distributions flatten the degradation curve.
3. **Monitor drift and outcomes** ($\mathcal{D}(P_t \| P_0)$). Divergence can trigger investigation; labeled outcomes or validated proxies determine whether retraining is warranted.

In practice, *knowing when to reevaluate is as important as knowing how to train*. A system can alert when $\mathcal{D}(P_t \| P_0) > \tau$ and retrain when outcome evidence confirms degradation. Without drift monitoring, it is blind to changing inputs. Chapter 14 develops the monitoring infrastructure and alerting strategies that implement this principle.

Population-dependent performance adds a failure mode that traditional software monitoring does not cover. Dependencies and environments can change traditional systems; ML models can also degrade because the input population changes. Monitoring must therefore extend beyond uptime and error rates to input distributions and labeled outcomes. Because exhaustive testing is impossible, continuous performance evaluation becomes an architectural requirement.


The same pattern appears in a lower-stakes setting. A product recommender trained on last season's click history might lose several percentage points under mild seasonal drift or tens of points under a severe training-serving skew, with the rate depending on the measured distribution shift and the model's sensitivity to that shift. This degradation often stems from **training-serving skew**, where features computed differently between training and serving pipelines cause model performance to degrade despite unchanged code. This is a systems issue that manifests as model-quality degradation.

These failure modes reshape the entire lifecycle. Teams must monitor infrastructure health, model quality, data quality, and prediction distributions together, then update the model when evidence confirms that production behavior has shifted. Silent degradation therefore has to inform system design from data collection through inference serving.

Silent degradation addresses the first half of the dual mandate by asking whether learned behavior remains trustworthy as the world changes. The second half asks whether the machine can produce that behavior within its time, memory, energy, and cost budgets. The bitter lesson established that computational scale drives AI progress; the question now becomes how to reason quantitatively about the data movement, computation, and overhead that constitute that scale.

1.7 Iron Law of ML Systems

The physical cost of computational scale becomes concrete in two familiar failures. A training job stalls when storage cannot feed an accelerator; an inference path misses its deadline when model state moves too slowly through memory or across the network. Their symptoms differ, but both consume the same finite time budget through data movement, computation, and fixed overhead.



The **iron law of ML systems** makes this shared structure explicit by decomposing total execution time T (seconds) into those three physical costs (equation 1.4):

$$T = \underbrace{\frac{D_{\text{vol}}}{\text{BW}}}_{\text{The Data Term}} + \underbrace{\frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}}_{\text{The Compute Term}} + \underbrace{L_{\text{lat}}}_{\text{The Latency Term}} \quad (1.4)$$

This equation is the mathematical spine of this book. It decomposes the total time required for any ML task, whether training a model for weeks or serving an inference in milliseconds, into three terms that correspond directly to the physical constraints of the dual mandate introduced in section 1.1:

1. **The data term** (D_{vol}/BW) represents the physical cost of moving bits. D_{vol} is the volume of data moved (bytes), and BW is the memory or network bandwidth (bytes/s). Whether loading terabytes from cloud storage or fetching weights from high-bandwidth memory, performance is often limited by I/O physics. Part I develops this foundation.
2. **The compute term** ($O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$) represents the cost of arithmetic. O is the number of floating-point operations (FLOPs), R_{peak} is the hardware’s theoretical peak throughput (FLOP/s), and η_{hw} is dimensionless realized hardware utilization ($0 \leq \eta_{\text{hw}} \leq 1$). Parts II and III develop this term.
3. **The latency term** (L_{lat}) represents the irreducible “tax” of system orchestration, networking, and serialization (seconds). This fixed latency dominates in real-time deployment. Part IV develops this term.

🔍 Systems Perspective 1.3: The iron law analogy

The name “iron law” follows Patterson & Hennessy’s Iron Law of Processor Performance (Hennessy and Patterson 2017). P&H’s law is a *multiplicative decomposition* (a tautology factoring CPU time), whereas this equation is an *additive first-order model* that approximates performance under simplifying assumptions.

The additive form assumes sequential execution. When movement and computation overlap, we replace the sum with their critical-path lower bound in equation 1.5:

$$T_{\text{pipelined}} \geq \max \left(\frac{D_{\text{vol}}}{\text{BW}}, \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} \right) + L_{\text{lat}} \quad (1.5)$$

This max-based formulation is the systems equivalent of overlapping asynchronous Direct Memory Access (DMA) data transfers with active Arithmetic Logic Unit (ALU) computation, where the execution time of the slower pipeline stage hides the latency of the faster stage. Equality requires ideal overlap and assumes that L_{lat} remains outside the overlapped phases. The term remains useful because, like Amdahl’s Law (Amdahl 1967), its value lies in identifying which physical constraint dominates before optimizing. The iron law simplifies the complexity of the full stack into three manageable terms. Section A.3.2 presents the refined treatment, including pipelining and overlap techniques that transform the additive model into the max-based formulation used in practice.

Once execution time has been decomposed, the next question is which term dominates. The Roofline Model answers the movement-versus-compute part of that question with a boundary called the **ridge point**. Below the ridge, data movement dominates; above it, arithmetic throughput dominates. Section D.2.1 develops the deeper mathematical treatment for readers who want it. Every optimization technique developed in this book manipulates one of these variables by moving less data, doing less work, using the machine more effectively, or reducing orchestration delay. A GPT-3-class training estimate makes that manipulation concrete by showing how an efficiency change propagates through the iron law.

Napkin Math 1.1: Training GPT-3

Problem: How long would it take to train a GPT-3-class model on a reference cluster with 1,024 accelerators?

Given:

- **Ops** (O): $\approx 3.14 \times 10^{23}$ FLOPs
- **Peak** (R_{peak}): 312 TFLOP/s
- **Efficiency** (η_{hw}): ≈ 45 percent (typical for large-scale distributed training)
- **Scale** (N_{accel}): 1,024 accelerators

Math:

$$T_{\text{train}} \approx \frac{O}{N_{\text{accel}} \cdot R_{\text{peak}} \cdot \eta_{\text{hw}}} \approx \frac{3.14 \times 10^{23}}{1024 \times 312 \times 10^{12} \times 0.45} \approx 25 \text{ days}$$

Result: 25 days.

Systems insight: If we improve hardware utilization (η_{hw}) from 45 percent to 60 percent through better scheduling and more efficient execution, training time drops to 19 days, saving 6 days of expensive compute time.


The equation is dimensionally consistent because each term resolves to seconds. One *cannot* add FLOPs to bytes any more than one can add meters to kilograms; the iron law adds time to time. A formal treatment in section D.2.2 verifies this consistency and demonstrates how unit tracking prevents common modeling errors.

The iron law governs *time*, but time is not the only constraint. For mobile devices, edge systems, and large-scale training clusters, *energy* often matters more than raw speed.

Just as time is governed by physics, so is energy. A fourth term enters the mental model through the **energy tax**. In many modern systems (mobile, edge, and large-scale training), energy, not time, is the hard constraint. Let D_{vol} be the total data volume moved (bytes), E_{move} the energy per byte moved, O the total operation count, and E_{compute} the energy per operation. Equation 1.6 formalizes this relationship, following the hardware-energy observation that data movement can dominate arithmetic energy (Horowitz 2014):

$$E_{\text{total}} \approx \underbrace{D_{\text{vol}} \times E_{\text{move}}}_{\text{Movement Term}} + \underbrace{O \times E_{\text{compute}}}_{\text{Compute Term}} \quad (1.6)$$

Per access, data movement can cost far more than arithmetic, so $E_{\text{move}} \gg E_{\text{compute}}$. Under the energy constants used in this text, moving one byte from off-chip Dynamic Random-Access Memory (DRAM) consumes roughly 145.5× the energy of an FP16 multiply and 800× that of an INT8 (8-bit integer) multiply (Horowitz 2014). Whether movement dominates a complete workload also depends on the number and width of transfers relative to its operation count. The physical reason for the per-access gap is that data movement requires charging and discharging wires over longer distances, while arithmetic occurs locally within a processing unit's circuits. Minimizing avoidable data movement (D_{vol}) can therefore improve both speed and energy efficiency.



Moving a byte costs about 145 times an FP16 op, the data-movement tax.

Checkpoint 1.2: The iron law

The iron law ($T \approx \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}}$) is the analytical backbone of this book. Before proceeding, verify you can manipulate its terms:

- ☐ **Data term** (D_{vol}/BW): identify the physical resource that bounds this term and name one workload regime where it dominates.
- ☐ **Compute term** ($O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$): identify the hardware quantity and utilization factor that bound this term.

- **Latency term** (L_{lat}): identify which costs remain after data movement and arithmetic have been optimized.
- **Peak throughput**: if you double R_{peak} , which term improves, and why do the other terms remain unchanged?

The same terms that determine time and energy also determine cost. Every byte moved, operation executed, and millisecond of latency consumes infrastructure budget. The next test is therefore economic, asking whether added compute buys enough model improvement to justify the resources it consumes.

1.7.1 Return on compute (RoC) as an economic lens

The decomposition of time also has an economic consequence. Following the quantitative reasoning tradition of Hennessy and Patterson, **return on compute (RoC)** measures the incremental accuracy gain per added dollar of infrastructure investment. With accuracy measured on a fixed scale, RoC has units of accuracy points per dollar.

$$\text{RoC} = \frac{\Delta \text{Accuracy}}{\Delta \text{Compute Cost}}$$

This ratio exposes an economic boundary. A 1-percentage-point gain in accuracy may fail the RoC test if it requires a $10\times$ increase in O (Total Operations). Every optimization in the following chapters targets either the numerator (extracting more signal from the same data) or the denominator (reducing the cost of executing the math). If the RoC is negative or negligible, the system is over-engineered, regardless of its technical sophistication. This economic lens transforms “accuracy” from a research target into an engineering budget.

If scale is the ultimate lever for performance, it is also the ultimate consumer of resources. The bitter lesson teaches that scale works, but the iron law teaches *how* to afford it. This tension between scaling and sustainability shapes the engineering principles that follow.

1.7.2 Lighthouse models put the iron law into practice

The iron law does more than diagnose bottlenecks; it organizes the entire discipline. Each term in the equation corresponds to a core engineering imperative. The data term demands *building* robust data pipelines and infrastructure (chapter 4). The compute term requires *optimizing* algorithms and hardware utilization for efficiency (Part III). The latency term necessitates *deploying* and *operating* systems reliably in production (chapter 13, chapter 14). These three imperatives structure the textbook. Parts I and II address building, Part III addresses optimization, and Part IV addresses deployment and operations.

Abstract equations become tangible through workloads. This textbook employs five recurring **lighthouse models** as diagnostic tools for the iron law. These canonical workloads reappear across chapters to test how the same physical constraints affect different architectural patterns.

Each lighthouse model represents a distinct stress case for the iron law. For instance, ResNet-50 provides a probe for compute throughput when a system repeatedly reuses the same learned parameters, while GPT-2/Llama acts as the primary probe for memory bandwidth pressure during language generation. For language models, autoregressive decode means generating one token at a time; the KV (Key-Value) cache is the saved attention state from previous tokens (chapter 6 introduces the attention mechanism in full), and prefill is the initial pass that processes the prompt before token-by-token generation begins. That diagnosis depends on the operating regime. Small-batch decode often streams weights and KV-cache state fast enough to expose memory bandwidth, while prefill and high-batch serving can shift the bottleneck toward arithmetic or communication. By following these same workloads from data engineering through to edge deployment, each chapter demonstrates how a single architectural choice propagates physical and economic constraints across the entire system.

The iron law makes these differences precise. ResNet-50 applies the same small weight filters across many spatial positions and, under batching, across many inputs; that reuse can make $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$

the dominant term because the processor must sustain enormous arithmetic throughput while the data footprint remains modest. GPT-2, by contrast, loads billions of unique weight parameters for every token it generates, and each weight is used only once before the next must be fetched; its D_{vol}/BW term dominates because memory bandwidth, not arithmetic, is the binding constraint. Applying the same equation to two different workloads leads to different diagnoses and therefore different optimization strategies. Doubling R_{peak} helps batched ResNet-50 once reuse increases computation per byte moved, but barely affects GPT-2 decode; doubling BW has the reverse effect for bandwidth-bound decode. Table 1.5 summarizes why each lighthouse model serves as a diagnostic tool for a specific bottleneck.

Table 1.5: Lighthouse Models as Reference Workloads: Each workload isolates a distinct bottleneck, enabling systematic investigation of how system constraints affect different architectural patterns. Quantitative specifications and architectural details appear in chapter 6.

Lighthouse Model	System Bottleneck	What It Reveals	Key Engineering Questions
ResNet-50	Compute throughput under reuse	GPU utilization, batching	Is the hardware doing math or waiting for data?
GPT-2/Llama	Memory bandwidth	Weight and sequence-state movement	How fast can model state move to compute?
Deep Learning Recommendation Model (DLRM)	Memory capacity	Embedding tables, scale-out	How do terabyte-scale models fit in memory?
MobileNetV2	Latency and power	Efficient operator design	Can the system meet real-time constraints on battery?
Keyword spotting	Power envelope	Tiny memory and energy budgets	Can the system run always-on inference on milliwatts?

Each lighthouse model manifests different constraints along the D-A-M axes, ensuring that the principles developed throughout this text are tested against the diversity of real-world systems engineering challenges. The division of labor among the book’s recurring examples is deliberate. The four deployment paradigms fix the *envelope* a system must operate within, the five lighthouse models supply the *workloads* that stress it, and in section 1.9.3 four engineering missions and three production case studies (Waymo, FarmBeats, and AlphaFold) pair envelope with workload under real-world constraints.

The same diagnostic reading applies retrospectively to the breakthrough that launched the deep learning era. The AlexNet system combined a convolutional architecture whose parallel matrix operations matched GPU capabilities with the 1.3M labeled images in the 2012 ImageNet challenge training split²¹ (Deng et al. 2009). Its error reduction therefore reflected coordination across the D-A-M axes rather than algorithmic novelty in isolation.

This interdependence means that optimizing one component often shifts pressure to another. AlexNet’s co-design success came at a cost affordable in 2012 (two consumer GPUs for a week), but modern models demand resources roughly 7 orders of magnitude larger. If the iron law governs *how fast* a system runs, a framework remains necessary to reason about *how efficiently* it uses those resources.

1.8 Three Dimensions of ML Efficiency

That efficiency question exposes a tension first introduced by the bitter lesson in section 1.4. Scale drives AI progress, but ever-larger datasets and compute budgets narrow participation to the most resource-rich organizations. Even those organizations eventually meet physical limits in data center power, memory bandwidth, and the diminishing returns of adding more parameters.

Common public estimates for GPT-4-class training place the compute budget around 2.5 million accelerator-days, representing millions of dollars in compute costs and substantial environmental impact. Many research institutions and companies cannot afford to compete through brute-force scaling. Those cost and access constraints make efficient use of existing compute a complementary path to progress.

Efficiency is a bottleneck diagnosis, not a single technique. The D-A-M landscape (figure 1.4) now becomes an action map. Data selection improves *what the system can learn from*, algorithmic efficiency reduces the work required to learn or predict, and compute efficiency aligns that work with the



GPT-2 decode is bandwidth-bound because the data term dominates.

²¹ | **ImageNet:** The 2009 paper reported 3.2 million images across 5,247 synsets; the later full dataset grew to about 14.2 million images across 21,841 synsets. The 2012 challenge training split used by AlexNet contained about 1.3M labeled images (Deng et al. 2009; Russakovsky et al. 2015) (see chapter 4).

machine. The remaining intersection, *how information moves*, cuts across all three because every improvement must survive the memory and communication path.

Algorithmic efficiency, the earliest frontier, reduces computational requirements through better model design and training procedures. Its goal is to produce more useful behavior per operation, so capability rises without scaling every resource in lockstep. As algorithms demanded ever more computation, **compute efficiency** became the second critical dimension. It maximizes hardware utilization by aligning algorithmic logic with machine physics, turning theoretical processor capability into useful work. Most recently, **data selection** emerged as the third dimension, extracting more learning signal from limited examples and thereby reducing the total operations term O of the iron law. The timeline in figure 1.5 places these three dimensions side by side before the chapter sequence presents them in build order. Together, these three dimensions provide the engineering tools to overcome the data, algorithm, and machine walls that pure scaling alone cannot address.




Figure 1.5: Historical Efficiency Trends: Three color-coded sequences align the algorithmic, compute, and data-selection eras. The paired dates beneath each box mark its interval: 1980–2010, 2010–2022, and 2023–future.

These three dimensions did not emerge simultaneously; each progressed through distinct eras at different rates. Algorithmic efficiency led the way, compute efficiency followed as demand grew, and data-centric methods matured most recently. While history progressed from algorithmic breakthroughs to hardware acceleration to data-centric methods, Part III reverses that sequence. Data selection comes first, followed by model compression and hardware acceleration. This pedagogical order reflects how practitioners build systems because quality data is prerequisite to effective model optimization, and understanding the model is prerequisite to mapping it efficiently onto hardware.

The trajectory of model architectures over time illustrates how these dimensions progress. Figure 1.6 makes the impact of algorithmic efficiency visible model by model, demonstrating how identical accuracy targets require progressively less compute as architectures improve.

The magnitude of efficiency improvements is measurable. Between 2012 and 2019, computational resources needed to train a neural network to achieve AlexNet-level performance on ImageNet classification decreased by approximately 44.5× (Hernandez and Brown 2020). This improvement, which halved about every 15 months, outpaced hardware efficiency gains predicted by Moore’s Law,²² demonstrating that algorithmic innovation drives efficiency as much as hardware advances.

Simultaneously, aggregate training compute in published frontier runs followed a much steeper cadence than Moore’s Law, with a fitted doubling time of approximately 3.4 months (Amodei and Hernandez 2018b). That aggregate publication trend is not the same quantity as an endpoint ratio between two landmark models, but it explains *why* efficiency optimization is not optional. Without it, only the most resource-rich organizations could participate in AI development.

These measurements emerge from rigorous empirical methodology that tracked training compute across hundreds of published models; chapter 12 develops the measurement frameworks that enable such systematic analysis of ML system performance. The two cadences just compared define the systems gap as the widening distance between what models demand (compute doubling every 3.4 months) and what hardware supplies (transistor density doubling roughly every two years). Closing that gap is the primary objective of this textbook, requiring integrated expertise across the software and hardware stack; chapter 11 quantifies it directly.

Architecture-by-architecture gains tell only half the story. Algorithmic improvements could not contain rising training demand. Figure 1.7 compares illustrative AlexNet-era and GPT-4-class

22 | Moore’s law: Gordon Moore’s 1965 observation described rapid growth in the number of components that could be economically integrated on a chip (Moore 1998); later industry summaries often expressed the cadence as roughly a two-year doubling.




Figure 1.6: Algorithmic Efficiency Trajectory: Training efficiency factor relative to AlexNet (2012 baseline) for ImageNet classification. Most later architectures achieve comparable accuracy with fewer computational resources, although individual points vary. The trajectory from AlexNet ($1\times$) through VGG (Visual Geometry Group), ResNet, MobileNet, and ShuffleNet to EfficientNet ($44\times$) demonstrates a 44.5-fold reduction in required compute over seven years, independent of hardware improvements (Hernandez and Brown 2020). Early-2020s points and the dashed fit through all plotted points are author-added context rather than evidence for extending the Hernandez-Brown trend.

endpoints, distinct from the aggregate 3.4-month trend. Their several-order-of-magnitude gap makes efficiency optimization necessary.




Figure 1.7: The Era of Scale: Illustrative training-compute estimates (FLOPs) vs. year on a log scale. While early deep learning (blue) showed rapid growth, the transformer era (red) accelerated this trend significantly. From AlexNet (2012) to illustrative GPT-4-class training scale anchors (2023), compute requirements increased by roughly 7 orders of magnitude ($16.7M$ times), far outpacing Moore's Law. This endpoint comparison is distinct from aggregate training-compute trend studies (Amodei and Hernandez 2018b; Sevilla et al. 2022) and is used here to motivate the specialized infrastructure described in this book.

Taken together, these two figures reveal a seeming contradiction that defines the economics of modern AI development. Figure 1.6 shows efficiency improving $44.5\times$ while figure 1.7 shows compute demand growing by roughly 7 orders of magnitude. The plots, however, hold different quantities constant. The first asks how much computation is required to reach a roughly fixed capability target, whereas the second permits the target to expand and records the resulting training budget. Efficiency lowers the cost of a given capability; scale determines how the newly affordable computation is spent. The apparent contradiction is therefore an economic feedback rather than a disagreement between the measurements.



Systems Perspective 1.4: The efficiency paradox

This dynamic is the ML systems equivalent of **Jevons paradox**. In resource economics, increasing the efficiency with which a resource is used often increases rather than decreases total consumption. In AI, efficiency gains fund larger models and broader datasets rather than smaller compute budgets. If EfficientNet needs 44.5× less compute than AlexNet to reach the same accuracy, organizations can invest the savings in larger models and broader data. This is precisely how GPT-3 came to require orders of magnitude more compute than AlexNet despite per-FLOP efficiency gains. The feedback loop, in which efficiency enables scale and scale demands efficiency, defines the modern AI engineering landscape.

The specific methods for achieving these gains are developed systematically in chapter 10 (algorithmic techniques) and chapter 11 (hardware foundations). Chapter 9 addresses data selection as an efficiency technique, while chapter 4 covers the pipeline design and quality infrastructure that make selected data usable.

1.9 AI Engineering as a Discipline

A cloud service may optimize throughput, while an edge device must remain within a strict power envelope. The same model can therefore be efficient in one setting and unusable in another. Model accuracy alone does not specify a working system: learned behavior must remain trustworthy as data changes, and the machine must deliver that behavior within its operating budget.

The degradation equation, iron law, and efficiency framework supply quantitative tools for this dual mandate. Together they span statistical behavior, computation, and deployment constraints. Applying them crosses disciplinary boundaries. Computer science addresses algorithms, and electrical engineering addresses hardware, but neither alone encompasses the integrated problem of building systems that remain reliable, efficient, and scalable in production. That problem defines AI engineering.



Definition 1.3: AI engineering

AI engineering is the discipline of designing, deploying, and maintaining ML systems that hold statistically evaluated behavior to deterministic reliability targets while satisfying production constraints across all three D-A-M axes: Data quality, Algorithm correctness, and Machine efficiency.

1. **Significance:** ML research typically optimizes only the algorithm axis (O and convergence). AI engineering jointly optimizes all three by bounding D_{vol} through data governance requirements, $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$ through production latency requirements, and total power draw through energy and cost budgets. A production system that achieves 95 percent accuracy in research but violates a 100 ms latency requirement in production is a failed system, regardless of its algorithm score.
2. **Distinction:** Unlike machine learning research, which targets a single objective (validation loss) on a static dataset, AI engineering targets a multi-objective constraint surface (latency, throughput, accuracy, cost, fairness, and robustness) on a distribution that shifts continuously after deployment.
3. **Common pitfall:** A frequent misconception is that AI engineering is just “software engineering for ML.” The system specification is instead probabilistic. An ML system’s output is statistically valid or invalid relative to a shifting distribution, not correct or incorrect relative to a fixed deterministic contract. This makes continuous monitoring a structural requirement, not an operational choice.

The phrase “stochastic systems with deterministic reliability” connects AI engineering to an earlier disciplinary convergence. Computer engineering emerged in the late 1960s and early 1970s²³ as computing systems grew too complex for electrical engineering or computer science to address

²³ | **Computer engineering:** Formalized as an academic discipline when Case Western Reserve launched the first accredited program in 1971, recognizing that neither electrical engineering nor computer science alone could address building reliable computers from unreliable components. ML systems engineering recapitulates this convergence: the binding constraint is not algorithmic or hardware in isolation but the integration of both under latency, power, and data-quality budgets that neither discipline’s curriculum addresses.

alone. It bridged both fields around the integrated problem of building reliable computers. AI engineering faces a similar problem at the intersection of algorithms, infrastructure, and operational practices.

AI engineering names the broader discipline; throughout this text, “ML systems engineering” names the practical work of designing, deploying, and maintaining modern ML systems.

1.9.1 Engineering across the ML lifecycle

A breakthrough algorithm requires efficient data collection and processing, distributed computation across hundreds or thousands of machines, reliable service to users with strict latency requirements, and continuous monitoring based on real-world performance. These obligations form a recurring lifecycle rather than a sequence that ends at deployment. The engineering object is no longer only code; it is code, data, model behavior, deployment context, and monitoring evidence evolving together. Production feedback can force a deployed system back into data collection and training, bending the familiar linear arc into a cycle.

The structural difference shows up first in tooling. Decades of established practice support code-defined behavior through version control that maintains precise histories, continuous integration pipelines that automate testing, and static analysis tools that measure quality. Behavior learned from data slips through this tooling, because the artifact that changes is no longer a diff a developer wrote. Chapter 3 develops the specialized workflows these challenges demand.

The deeper difference is the prominence of continuous feedback. The loops in figure 1.8 show why. When monitoring detects performance degradation, the system does not receive a code patch alone. It may cycle back through data collection, preparation, training, and evaluation before redeployment, making iteration part of the operating architecture rather than only the development process.




Figure 1.8: ML System Lifecycle: A six-box flowchart depicting Data Collection, Preparation, Model Training, Evaluation, Deployment, and Monitoring. Two feedback loops distinguish this cycle from linear software development: evaluation returns to preparation when results are insufficient, and monitoring triggers new data collection when performance degrades.

The data-dependent nature of ML systems creates dynamic lifecycles requiring continuous monitoring and adaptation. Unlike source code that changes only through developer modifications, data reflects real-world dynamics, and distribution shifts can silently alter system behavior without any code changes. The tooling gap identified earlier in this section follows the system into production. Version control built for discrete code changes struggles with large, evolving datasets, and testing frameworks built for deterministic outputs require adaptation for probabilistic predictions. Chapter 4 develops data versioning and quality management, while chapter 14 develops monitoring for probabilistic behavior.

What each lifecycle stage demands is not uniform; it depends on the mission the system is built to serve.

Systems Perspective 1.5: From paradigms to missions

The top of the hierarchy transforms abstract systems into concrete engineering missions. Each mission inherits one of the four deployment paradigms introduced in section 1.5.1 and pairs it with a scenario-specific workload. The four missions in table 1.6 are the application scenarios that recur throughout the book and its associated labs. They act as end-to-end tests of engineering decisions. A $2\times$ increase in memory bandwidth is an academic result until it is proven to extend the battery life of the Smart Doorbell or enable the safety-critical latency required for Autonomous Perception.

Each mission runs this lifecycle continuously. High-quality data can improve the model, which improves the product and the feedback it generates. Conversely, weakness at any stage can propagate through everything downstream.

Table 1.6: Four Engineering Missions: Each deployment paradigm is paired with a recurring scenario workload and its binding constraint.

Mission	Deployment Paradigm	Scenario Workload	Critical Constraint
Frontier training	Cloud Cluster	GPT-4	Target: 500 ms/step
Autonomous perception	Edge Robotics	YOLOv8-nano	SLA (Service-Level Agreement): 10 ms latency
Mobile assistant	Smartphone	Mobile-optimized small LLM (Large Language Model)	RAM: < 2 GB / Thermal: < 3 W
Smart Doorbell	TinyML (MCU [Microcontroller Unit])	Wake Vision	Power: 100 mW

1.9.2 Deployment context shapes the lifecycle

Deployment context determines which lifecycle pressures dominate. The same stages apply across ML systems, but a megawatt-scale data center and a milliwatt-scale embedded device impose different bottlenecks on data collection, model updates, monitoring, and serving.

At one end of the spectrum, cloud-based ML systems train large models and serve millions of users, trading abundant computing resources for capacity limits, operational complexity, and high costs. Chapter 2 examines their architectural patterns, while chapter 11 develops the hardware foundations that make this scale economically viable.

At the other end, TinyML systems run on microcontrollers²⁴ and embedded devices. Their kilobyte-scale memory and milliwatt power budgets make feasibility precede model quality: a smart-home device must recognize a command using less power than an LED bulb, while a sensor may need to detect anomalies on one battery for years. The efficiency framework in section 1.8 supplies the governing principles, while chapter 10 develops the techniques that make such deployment possible.

Between these poles, placement becomes a constraint-allocation problem. Edge ML systems move computation toward data sources to reduce latency²⁵ and bandwidth demand. Mobile ML systems share memory, thermal headroom, and battery power with every other application, trading raw speed for locality and privacy. Hybrid systems distribute work across tiers to balance latency, privacy, bandwidth, and update control.

Each position on this deployment spectrum creates distinct bottlenecks that determine which efficiency dimensions matter most, as summarized in table 1.7:

Table 1.7: Efficiency Priorities by Deployment Context: Each deployment environment creates distinct bottlenecks, requiring tailored optimization strategies. Cloud systems optimize for throughput and cost; edge systems optimize for memory and power; TinyML systems require extreme efficiency across all dimensions.

Environment	Primary Constraint	Efficiency Focus
Cloud training	Cost, throughput	Distributed efficiency, hardware utilization
Cloud inference	Latency, cost per query	Batching, model serving optimization
Edge devices	Memory, power	Smaller models and lower data movement
Mobile	Battery, thermal	Energy-efficient inference
TinyML	kilobyte-scale memory, mW power	Extreme compression, specialized architectures

The deployment spectrum represents more than different hardware configurations. Each deployment environment reshapes every stage of the ML lifecycle, from initial data collection through continuous operation and evolution, creating an interplay of constraints that traditional software rarely encounters.

Consider how a single deployment decision cascades through the entire system. Latency-sensitive applications like autonomous vehicles or real-time fraud detection require edge or embedded

²⁴ **Microcontrollers:** Single-chip computers with kilobytes of memory and milliwatts of power budget. For TinyML, memory and energy determine feasibility before model quality.

²⁵ **Latency:** From Latin *latere* (“to lie hidden”), delay is invisible until it causes failure. At 30 m/s, every millisecond adds 3 cm of travel before braking begins, making L_{lat} the edge constraint.

architectures despite their resource constraints, while large language models naturally gravitate toward centralized cloud infrastructure. This initial architectural choice, however, determines far more than where computation happens. Cloud systems must optimize for cost efficiency at scale, balancing expensive GPU clusters, storage, and network bandwidth, which in turn shapes how often models are retrained, what historical data is retained, and how inference load is distributed. Edge and mobile systems face fixed resource limits that constrain model complexity and update frequency, forcing aggressive model compression²⁶ and careful scheduling. The strictest constraints arise in embedded and TinyML environments, where every byte of memory and milliwatt of power matters.

Operational complexity increases as systems become more distributed. Centralized cloud architectures benefit from mature deployment tools and managed services, while edge and hybrid systems must coordinate data collection across sensors with varying connectivity, track models deployed across thousands of devices, handle staged rollouts with rollback capabilities, and aggregate monitoring signals from geographically distributed endpoints (chapter 14). Data considerations introduce competing pressures. Privacy requirements or data sovereignty regulations may push computation toward the edge, while the need for large-scale training data pulls toward centralized cloud aggregation. Model updates also behave differently across the spectrum. Cloud architectures enable rapid iteration through centralized traffic control, while edge deployments require remote updates with careful bandwidth management and rollback capabilities.

In practice, these trade-offs are rarely simple binary choices. Modern ML systems often adopt hybrid approaches that span the deployment spectrum. An autonomous vehicle performs real-time perception and control at the edge for latency reasons, uploads driving data to the cloud for model improvement, and periodically downloads updated models. A voice assistant runs wake-word detection on-device to preserve privacy and reduce latency but sends full speech to the cloud for complex natural language processing. The key insight is that a choice to deploy on embedded devices constrains more than model size; it affects data collection strategies, training approaches, evaluation metrics, deployment mechanisms, and monitoring capabilities. These interconnected decisions demonstrate the D·A·M taxonomy in practice, where constraints along one axis create cascading effects throughout the system.

Three production systems make these abstract trade-offs concrete by representing the extremes of the deployment spectrum. Each system faces the same core challenges (data quality, model complexity, and machine scale), but the constraints of its deployment environment force radically different engineering solutions.

1.9.3 Production systems expose shared challenges

A deployment case study becomes an engineering tool when it exposes the binding constraint behind a design. The three production case studies (Waymo, FarmBeats, and AlphaFold) sit at different extremes of the deployment spectrum, so the same D·A·M questions force different engineering responses.

- **Autonomous driving**²⁷ binds on safety-critical latency and data freshness. The Waymo Open Dataset provides camera and LiDAR data collected across varied driving environments (Sun et al. 2020), illustrating the multimodal inputs and geographic coverage that a perception stack must handle. The broader case uses a representative high-stakes hybrid pattern with on-vehicle inference for low latency and cloud infrastructure for training and evaluation.
- **FarmBeats**²⁸ (Vasisht et al. 2017) binds on connectivity and data freshness. Microsoft's precision agriculture platform connects field sensors to a local gateway PC for edge processing. TV white-space networking carries data across the farm, while the weaker farm-to-cloud Internet link constrains synchronization.
- **AlphaFold** (Jumper et al. 2021) binds on compute-intensive training and curated scientific data. DeepMind's protein structure prediction system made a landmark advance on a 50-year grand challenge in biology. AlphaFold represents the compute-intensive cloud deployment pattern. Initial training used 128 TPUs for approximately one week, followed by about four days of fine-tuning, and drew on the Protein Data Bank's experimentally determined structures.

26 | Model compression: A family of techniques, including quantization, pruning, and distillation, that reduces model storage or computation. Its size and accuracy effects depend on the model, method, workload, and target hardware.

27 | Autonomous-driving hybrid workflow: This representative workflow forces a synchronization challenge absent from pure cloud or pure edge systems. The on-vehicle model must be controlled and regression-tested before deployment, while cloud infrastructure can train and evaluate improved versions on newly collected driving data. This creates a version-management gap between deployed and newly trained models, requiring rigorous validation before any remote model update can be pushed to safety-critical vehicles.

28 | FarmBeats: The system uses TV white-space links as high-bandwidth intra-farm backhaul from sensors to a gateway PC, where local processing reduces dependence on the weaker Internet connection to the cloud. The resulting constraint is timely data synchronization across that farm-to-cloud link, not delivery of a particular model size (Vasisht et al. 2017).

These systems complement the lighthouse models by illustrating how the same core challenges (data quality, model complexity, and infrastructure scale) manifest under radically different constraints. Rather than examining each system in isolation, they are analyzed through the lens of the D·A·M taxonomy. The same data drift phenomenon that affects Waymo’s perception models in changing weather also affects FarmBeats’ crop disease detection across growing seasons, though the engineering responses differ based on machine constraints.

The interdependencies across the D·A·M axes create specific challenge categories that define the daily work of an ML systems engineer. Examining the deployment extremes reveals these challenges in their most rigorous forms.

Real-world data is often noisy and inconsistent, presenting the first category of challenges. Autonomous vehicles process large multimodal sensor streams from LiDAR²⁹ and cameras (Sun et al. 2020). Engineers must solve for sensor interference, such as rain obscuring cameras, and temporal misalignment across asynchronous data streams. Scale compounds these quality issues. FarmBeats processes sensor data at a local gateway before synchronizing over a constrained farm-to-cloud link, while AlphaFold occupies the opposite extreme, requiring access to the Protein Data Bank’s experimentally determined structures during training.

Data drift creates an ongoing operational burden atop both quality and scale. The statistical properties of input data change over time, and models are only as reliable as their alignment with the current distribution (Gama et al. 2014; Quiñonero-Candela et al. 2009; Koh et al. 2021). The Waymo Open Dataset reports pronounced domain gaps among San Francisco, Phoenix, and Mountain View (Sun et al. 2020);³⁰ detecting such regional shifts requires continuous monitoring of input statistics before they manifest as system failures.

Beyond data, model complexity and generalization form the second challenge category. Computational intensity defines the upper bound of capability. Foundation models at GPT-3 scale (section 1.3.3) demand zettaFLOPs of compute, and even smaller scientific models like AlphaFold required weeks of specialized accelerator training. Systems engineers must optimize for “FLOP/s per watt” to make these models economically and environmentally viable. Yet raw scale is not enough. The **generalization gap** remains the central algorithmic risk because a model might achieve 99 percent accuracy on benchmarks but only 75 percent in the real world. For Waymo’s safety-critical autonomous driving systems, minimizing this gap is a life-or-death requirement, demanding robustness methods that cover the long tail of edge cases.

The third category encompasses the system-level challenges of getting models to work reliably in production. The **training-serving divide** describes the gap between the flexible environment where models are born and the rigid environment where they operate. Latency-throughput trade-offs dictate architecture. Waymo-style perception systems require low-latency safety decisions at the edge, while AlphaFold runs in the cloud and its inference time depends on protein length and configuration. Hybrid coordination adds further complexity, as modern systems increasingly adopt tiered architectures. A voice assistant, for example, performs wake-word detection locally (TinyML) to preserve privacy and reduce latency, but offloads complex natural language processing to massive GPU clusters in the cloud.

Finally, as systems scale, their impact on society becomes a first-class engineering concern that cuts across all three D·A·M axes. Fairness and bias must be managed proactively, since models can unintentionally learn societal biases present in their training data. Responsible engineering requires systematic auditing of performance across demographic subgroups to ensure equitable outcomes. Transparency and privacy requirements further constrain design. Many deep networks function as “black boxes,” yet in domains like healthcare or finance, stakeholders require interpretability. Systems must also be resilient against inference attacks³¹ that attempt to extract sensitive training data from model predictions.

The failure chain can cross all four challenge categories, so no single specialty can own it end to end. Their overlap creates an organizational problem as much as a technical one.

1.10 Five-Pillar Framework

Production ML therefore needs a framework that assigns responsibility for data, model behavior, infrastructure, and ethics without severing their connections (Paley et al. 2022). Traditional

²⁹ | **LiDAR (light detection and ranging):** This sensor is a primary reason the vehicle is a “roving data center,” as its pulsed lasers generate a dense 3D point cloud of the environment. The raw data stream from a single unit can exceed 100 megabytes per second, creating both the terabyte-scale volume challenge and the quality challenge mentioned, as the signal is easily degraded by sensor interference from rain or fog.

³⁰ | **Data drift:** Divergence between the training data distribution (P_0) and the production distribution (P_t). Drift can change performance without a code change, but divergence alone does not determine whether accuracy falls; outcome monitoring is needed to establish degradation (see chapter 14).

³¹ | **Inference attack:** A security threat where an adversary queries a model to deduce sensitive information about the training set. These attacks exploit the tendency of overparameterized models to memorize unique patterns in their training data, creating a direct trade-off between model capacity and privacy risk that motivates defensive techniques such as differential privacy and output perturbation.




Figure 1.9: Five-Pillar Framework: The architectural foundation of ML systems engineering organized into five core disciplines: Data Engineering, Training Systems (Model Training), Deployment Infrastructure (Model Deployment), Operations & Monitoring (Operation & Maintenance), and Ethics & Governance, supported by foundational efficiency, evaluation, and reliability practices.

software engineering practices alone cannot address systems that degrade quietly rather than failing visibly.

This work organizes ML systems engineering around five interconnected disciplines that directly address these challenge categories. Figure 1.9 presents five engineering pillars, each targeting a distinct challenge category and resting on a shared foundation that reflects the physical and economic constraints every pillar must respect. Together, they represent the core engineering capabilities required to bridge the gap between research prototypes and production systems capable of operating reliably at scale. While these pillars organize the practice of ML engineering, they are supported by the foundational technical imperatives of Performance Optimization and Hardware Acceleration (covered in Part III), which provide the efficiency required to make large-scale training and deployment economically and physically viable.

The pillars are easiest to understand through a failure chain. Suppose a wake-word model stops working reliably for users in a noisy apartment building after a model update. The first question is whether the training data captured that acoustic environment, whether labels were reliable, and whether the pipeline can trace which examples reached the model. The **data engineering** pillar (chapter 4) owns the data-quality, scale, privacy, drift, and lineage problems that determine what the model can learn.

If the data is sound, the next question is whether the training process converted it into a model that fits the task and budget. The **training systems** pillar (chapter 8) owns that boundary by coordinating datasets, frameworks, optimization algorithms, hyperparameters, distributed jobs, restarts, and the cost-quality trade-offs created by model scale. A model that trains successfully is still not a system. The **deployment infrastructure** pillar owns the training-serving divide across model packaging, inference performance, latency, throughput, device constraints, and the benchmarking methods that reveal whether the deployed artifact still meets the requirement.

Once the model is serving, failure becomes temporal. The **operations and monitoring** pillar owns the question of whether behavior remains acceptable after launch, when data distributions shift, traffic changes, and model quality can degrade while infrastructure dashboards stay green. It connects monitoring, alerting, rollout strategy, incident response, and continuous evaluation. Finally, the wake-word failure may not affect all users equally, and the audio pipeline may raise consent or privacy obligations. The **ethics and governance** pillar (chapter 15) owns constraints spanning fairness, transparency, privacy, safety, documentation, and accountability throughout the lifecycle.

Alternative organizational frameworks could group these concerns by component or lifecycle phase. The five-pillar structure was chosen because it matches the ownership boundaries that appear in real engineering teams while still making their interdependence explicit. Data choices shape training outcomes; training choices constrain deployment; deployment choices determine what

operations can observe; and governance requirements can change all four. Treating responsible engineering as its own pillar prevents it from becoming an implicit afterthought under deadline pressure.

Together, the pillars translate the D-A-M taxonomy (section 1.5.1) and lifecycle stages (section 1.9.1) into engineering ownership. This structure reflects how AI evolved from algorithm-centric research to systems-centric engineering, shifting focus from making individual algorithms work to building systems that reliably deploy, operate, and maintain those algorithms at scale.

These pillars also provide the organizational backbone for this textbook. Each part develops the knowledge and skills needed for one or more pillars, following the same progression engineers use to build systems in practice. Foundations come first, followed by model construction, optimization, and production deployment.

1.11 Book Organization

The five pillars describe what ML systems engineers must coordinate; the book's four parts describe the order in which readers develop those capabilities. The organizing principle is *context before theory*. The landscape and vocabulary are established (Part I) before building models (Part II), optimizing those models (Part III), and deploying them reliably (Part IV). Table 1.8 outlines this progression.

Table 1.8: Book Organization: The four parts follow a pedagogical progression from context (Foundations) through theory (Build) to practice (Optimize, Deploy). Each part builds on the vocabulary and frameworks of its predecessors, so Part III's optimization techniques assume familiarity with Part II's model architectures, and Part IV's deployment practices assume mastery of Parts II and III.

Part	Theme	Key Chapters
I: Foundations	Context: ML systems landscape	This chapter, chapter 2, chapter 3, chapter 4
II: Build	Theory: Model fundamentals	chapter 5, chapter 6, chapter 7, chapter 8
III: Optimize	Efficiency: Performance tuning	chapter 9, chapter 10, chapter 11, chapter 12
IV: Deploy	Production: Real-world systems	chapter 13, chapter 14, chapter 15, chapter 16

Part I establishes the constraint vocabulary before model machinery appears. This opening chapter develops the engineering revolution in AI and the frameworks that organize this discipline. Chapter 2 explores the deployment spectrum from Cloud to TinyML, examining how physical constraints (power envelopes, memory hierarchies, and latency budgets) govern each tier. Chapter 3 presents the end-to-end process from problem formulation through deployment, providing the conceptual map that guides subsequent learning. Chapter 4 addresses data collection, processing, and management, establishing that data infrastructure precedes and enables model development.

Part II turns that vocabulary into model construction skills. Chapter 5 provides algorithmic foundations, while chapter 6 extends these to specific network designs. Both chapters reference the five lighthouse models introduced in section 1.7.2 (ResNet-50, GPT-2/Llama, MobileNetV2, DLRM, and Keyword Spotting) to anchor abstract concepts in concrete workloads. Chapter 7 examines the software infrastructure from TensorFlow and PyTorch to specialized tools. Chapter 8 develops training systems for complex models and large datasets.

Part III asks how to change the terms of the iron law without losing quality. Chapter 9 introduces techniques for reducing computational requirements while maintaining quality. Chapter 10 covers model-reduction techniques that make deployment cheaper. Chapter 11 covers GPUs and application-specific integrated circuits (ASICs). Chapter 12 establishes methodologies for measuring and comparing system performance.

Part IV returns optimized systems to production, where degradation and deployment context dominate. Chapter 13 covers infrastructure for delivering predictions with low latency. Chapter 14 encompasses practices from monitoring and deployment to incident response. Chapter 15 addresses ethical considerations and governance. Chapter 16 synthesizes the complete methodology and prepares the reader for the transition from single-node mastery to fleet-scale orchestration.

This book covers the **single-node** regime of one host with one to eight accelerators, each typically using local device memory and communicating through an on-node interconnect. The binding

constraint depends on the workload and may be device memory capacity, memory bandwidth, computation, or interconnect communication. At fleet scale, thousands of nodes coordinate across network fabrics and the bottleneck shifts toward bisection bandwidth, the aggregate capacity across a cut through the cluster network. For detailed guidance on reading paths, learning outcomes, prerequisites, and how to get the most from this textbook, the preface provides the orientation.

The frameworks introduced in section 1.8 and section 1.10 help only if practitioners also shed assumptions carried over from adjacent fields. Every discipline accumulates intuitions that work within its boundaries but fail when applied elsewhere. ML systems engineering is particularly vulnerable to such imported assumptions because it draws from software engineering, statistics, and hardware design simultaneously, each of which cultivates subtly different intuitions about how systems should behave.

1.12 Fallacies and Pitfalls

Assumptions that hold in traditional software, academic research, or pure mathematics fail when applied to systems whose behavior emerges from data. The following fallacies and pitfalls capture errors that waste engineering effort, delay deployments, and cause silent production failures.

Fallacy: *Better algorithms automatically produce better systems.*

Engineers assume algorithmic sophistication drives system performance, but this ignores the iron law (section 1.7). Vision transformers demonstrate that architecture and large-scale pretraining can produce strong image-recognition results (Dosovitskiy et al. 2021), but production utility still depends on compute, memory movement, and latency budgets. In production, a model that is 1 percent more accurate but violates latency requirements has effectively zero utility. Production model selection is therefore a constrained optimization problem: maximize task quality subject to latency, memory, energy, cost, and reliability budgets. The hidden technical debt surrounding production models shows why model code is only the visible center of a much larger system. A well-engineered system with a simpler model can outperform a more sophisticated architecture lacking robust infrastructure.

Pitfall: *Treating ML systems as traditional software that happens to include a model.*

Engineers apply traditional testing and deployment practices to ML systems, but these systems fail in qualitatively different ways (section 1.6). Traditional bugs often produce immediate failures; ML systems can silently degrade over weeks or months before anyone notices. A/B tests in conventional software may show clear signals quickly, while ML comparisons can require longer observation windows to detect small accuracy differences across subpopulations. Unit tests verify deterministic paths; ML systems require monitoring infrastructure to catch unreliable predictions, data drift, and calibration failures. Teams deploying ML with only Continuous Integration and Continuous Delivery (CI/CD) pipelines risk silent failures that surface only after user-facing behavior has already degraded.

Fallacy: *High accuracy on benchmark datasets indicates production readiness.*

Engineers assume benchmark performance predicts production accuracy, but distribution shift and operational differences can cause substantial degradation in deployment. A sentiment analysis model that performs well on curated test data may fail sharply in production as users employ slang, emojis, and context absent from benchmarks. The deployment spectrum (section 1.9.2) shows that cloud, edge, and mobile environments each introduce distinct constraints. Network latency adds overhead, mobile devices' limited numerical precision can alter accuracy, and edge devices may lack the memory for multi-model strategies that boosted benchmark scores. Production systems require failure mode analysis across demographic subgroups, monitoring infrastructure to detect drift, and validation protocols that match actual operating conditions rather than idealized test sets.

Pitfall: *Optimizing individual components without considering system interactions.*

Engineers optimize inference latency in isolation, but Amdahl's Law governs end-to-end performance. A team reduces model inference from 45 ms to 15 ms, expecting proportional improvement. Yet preprocessing consumes 60 ms and postprocessing adds 25 ms, so total latency drops only from 130 ms to 100 ms. That is a 23 percent improvement rather than the expected 67 percent. The

D·A·M landscape (figure 1.4) shows that the Data, Algorithm, and Machine axes form an interdependent system where optimizing one component shifts bottlenecks rather than eliminating them. Component-level gains do not determine end-to-end improvement; the result depends on how much of the full path the optimized component occupies.

Fallacy: *ML systems can be deployed once and left to run indefinitely.*

Engineers assume deployed systems maintain performance indefinitely, but distribution shift can change a fixed model's performance. In this illustrative scenario, a recommendation system deployed at 85 percent accuracy drops to 80.2 percent within 6 months as purchasing patterns shift, losing 4.8 percentage points without any code changes. The ML lifecycle (section 1.9.1) therefore treats outcome monitoring and evidence-based retraining as operational requirements. Fraud-detection and Natural Language Processing (NLP) systems face the same risk because attackers adapt, vocabulary shifts, and user behavior changes while the code remains unchanged. Without monitoring, systems can appear healthy while prediction quality erodes. Organizations treating deployment as one-time often discover failures only after customer complaints or downstream metrics reveal the degradation.

Pitfall: *Assuming that ML expertise alone is sufficient for ML systems engineering.*

Organizations hire ML researchers expecting production-ready systems, but the five-pillar framework (section 1.10) requires integrated expertise across algorithms, software, systems, and operations. Teams with strong ML skills but limited systems experience can miss throughput targets because Application Programming Interface (API) design, storage layout, and serving infrastructure shape realized performance. Conversely, software infrastructure built without ML awareness can introduce preprocessing or feature bugs that degrade model behavior without obvious system failures. Deployment case studies show that production ML requires coordinated attention to data, models, infrastructure, and organizational workflow, not algorithmic quality alone (Paleyes et al. 2022). Effective teams integrate ML researchers, software engineers, and operations specialists rather than expecting one role to master all skills. The summary returns these failures to the chapter's central claim that ML systems engineering exists because learned behavior, physical infrastructure, and organizational workflow must be designed together.

1.13 Summary

Machine learning systems must satisfy two obligations at once: learned behavior must remain trustworthy, and the machine must deliver that behavior within physical and economic limits. The Software 2.0 shift explains why behavior learned from data can fail silently as distributions change. AI's paradigm history and the bitter lesson explain why progress repeatedly came from systems that could exploit more computation rather than from hand-coded expertise. The D·A·M taxonomy locates the binding constraint, while the degradation equation, iron law, and energy and efficiency frameworks turn those constraints into quantitative diagnoses.


The lifecycle, deployment spectrum, and production case studies then show why continuous iteration and context-aware design are mandatory. Five lighthouse models (ResNet-50, GPT-2/Llama, MobileNetV2, DLRM, and Keyword Spotting, detailed in chapter 6) recur throughout the book to ground these principles in real workloads.

Return to the smartphone interaction that opened the chapter. What appeared to be one intelligent action depended on representative data, a learned model, a machine within its operating budget, and feedback that could reveal changing behavior. That chain answers the question posed at the outset. Because ML behavior is learned as well as coded, it can degrade without an explicit failure and must be co-designed across data, algorithms, software, and hardware. AI engineering holds that stochastic behavior to deterministic reliability targets.



Key Takeaways: Constraints drive architecture

- **D·A·M bottlenecks migrate rather than disappear.** Data, Algorithm, and Machine constraints interact, so improving one axis often exposes another. The systems habit is to ask which axis now binds, then choose the intervention that relieves that constraint without creating a larger downstream failure.



- **Learned behavior can decay silently.** Traditional software usually fails when code or its environment changes; ML systems can degrade while code and infrastructure stay fixed because the world shifts relative to the training distribution. Drift metrics turn that shift into investigation triggers rather than surprise accuracy loss.
- **The iron law makes latency diagnostic.** Data movement, computation, and overhead all spend from the same time budget. Cutting inference from 45 ms to 15 ms gives only 23 percent improvement when preprocessing (60 ms) and postprocessing (25 ms) dominate, so optimize the term that binds end-to-end behavior.
- **Scale wins inside physical limits.** The bitter lesson explains why general methods with more compute displaced hand-crafted systems, but scale only helps when data, architecture, and machine can support it. Efficiency gains of 44.5× coexisted with roughly 7 orders of compute growth.
- **AI engineering is continuous co-design.** Deployment context, lifecycle monitoring, and the five engineering pillars are not later add-ons; they are how stochastic learned behavior is held to deterministic reliability targets from cloud training through TinyML operation.

Everything this chapter has introduced supports one claim: a machine learning system is governed by physics, not by intention. Its behavior reflects what its data, arithmetic, and hardware permit. The bitter lesson, iron law, degradation equation, and D·A·M taxonomy form a single vocabulary for reasoning about behavior that is learned rather than fully specified and can decay unless maintained. Treating those constraints as the real specification is what turns a collection of techniques into a discipline.



What's Next: From vision to architecture

Selecting where an ML model should run is governed by physical laws. The speed of light makes distant cloud servers useless for emergency braking. Thermodynamics prevents data-center-class models from running on a mobile device. Memory physics creates bandwidth ceilings that faster chips cannot overcome. The four deployment paradigms are where those laws land. Chapter 2 derives each paradigm's operating envelope from the physics and develops the decision framework for choosing among them when requirements conflict.

I

FOUNDATIONS OF ML SYSTEMS

Part I

Foundation Principles

Machine learning systems exhibit a deceptively simple heuristic: complexity often moves rather than disappearing, although good design can remove accidental complexity. Complexity flows among the three domains of the D·A·M taxonomy: **Data** as information, **Algorithm** as logic, and **Machine** as physics. Simplifying one domain can burden the others. A hand-crafted feature pipeline reduces algorithmic complexity but demands more data engineering effort. A larger model absorbs messy data but shifts complexity onto the hardware that must train and serve it. This **Conservation of Complexity** heuristic motivates everything in this book. The quantitative bounds, models, and principles introduced throughout the book describe constraints that emerge from where complexity currently resides.

Architectures, frameworks, and optimizations succeed only when they respect constraints imposed by hardware, mathematics, and information theory. Just as civil engineers cannot ignore gravity, ML engineers cannot ignore the physical laws that govern data, computation, and system throughput. Part I establishes these foundations: not best practices that evolve with frameworks or opinions that differ between teams, but enduring physical constraints and quantitative models for ML engineering. The first constraint starts with data itself, where the familiar boundary between program and input begins to disappear.

Principle 1: The data-as-code principle

Invariant: Data functions as the source code of the ML system. A change to the training dataset can alter the system’s learned behavior. Here, learned behavior means the trained input-output mapping; f denotes the training process, and $\approx$ marks a qualitative dependency rather than a predictive equality. Model architecture, initialization, optimizer configuration, and training stochasticity are included in the training procedure.

$$\text{Learned Behavior} \approx f(\text{Training Data, Training Procedure})$$

Implication: Data engineering requires the same rigor as software engineering. Datasets should be **versioned** (like Git), **unit-tested** (data quality checks), and **debugged**. Deleting a row of training data can alter the rebuilt model just as deleting a line of code can alter a compiled artifact.

If data functions as source code, then it is not merely a logical artifact; it also has physical properties that constrain system architecture when data and computation occupy different locations.

Principle 2: The data-gravity principle


Invariant: As data volume (D_{vol}) grows, transferring the dataset can cost more than sending a much smaller query or code payload over the same constrained link (McCrory 2010). Here, cost denotes one transfer metric at a time, such as elapsed transfer time or bytes sent, under identical link conditions; it does not combine latency and bandwidth into one quantity.

$$\text{Cost to transfer dataset} \gg \text{Cost to transfer query or code} \quad \text{in this regime}$$

Implication: Large datasets can become the gravitational center of the architecture. Systems may move compute to data by shipping queries or code to the storage layer rather than moving data to compute by repeatedly downloading large datasets.

Together, these two principles establish that data can be both the logical program and a physical anchor of an ML system. Part I develops these implications as a dependency chain. Deployment constraints define where a system can execute; the lifecycle carries those constraints through problem definition, data, modeling, validation, and operation; and data engineering makes them enforceable in versioned, tested pipeline artifacts. Ignoring any link returns complexity as integration cost: an accurate model that cannot execute, a workflow that cannot reproduce its inputs, or a pipeline whose outputs no longer match deployment conditions.

ML Systems




2.1	Deployment Paradigm Framework
2.2	Physical Constraints: Why Paradigms Exist
2.3	Analyzing Workloads
2.4	System Balance and Hardware
2.5	Cloud ML: Computational Power
2.6	Edge ML: Latency and Privacy
2.7	Mobile ML: Offline Intelligence
2.8	TinyML: Ubiquitous Sensing
2.9	Paradigm Selection
2.10	Hybrid Architectures
2.11	System Entropy: Why Deployment Is Not the End
2.12	Fallacies and Pitfalls
2.13	Summary

Purpose

Why does deploying the same model to a phone vs. a data center demand fundamentally different engineering?

The defining insight of ML systems engineering is that constraints drive architecture. The speed of light sets an absolute floor on how quickly distant servers can respond. Thermodynamics limits how much computation can occur in a given volume before heat becomes unmanageable. Memory physics makes moving data often more expensive than processing it. These are not engineering limitations awaiting better technology; they are permanent physical boundaries that partition the world into fundamentally distinct operating regimes. A data center can train billion-parameter models but cannot guarantee low-latency responses to users thousands of miles away. A smartphone can respond instantly but has a fraction of the memory budget. A microcontroller can run on a coin-cell battery for years but has barely enough compute for a simple keyword detector. The same model, the same algorithm applied to the same data, demands radically different engineering in each regime, not because the physics changes but because different constraints dominate. Teams that treat deployment as an afterthought, training in one environment and deferring target constraints until launch, discover too late that the target environment invalidates months of architectural decisions. The engineering task is therefore not to identify a universally best deployment but to locate the constraint that binds in the intended operating regime. In D-A-M terms, understanding these regimes transforms deployment from an operational detail into a first-order co-design problem in which data locality, algorithm structure, and machine constraints jointly determine what is possible.





Learning Objectives

- Explain how physical constraints create deployment paradigms from cloud to TinyML
- Apply the iron law and bottleneck principle to classify compute-, memory-, and I/O-bound workloads
- Map workload archetypes to deployment paradigms using lighthouse model examples
- Compare cloud, edge, mobile, and TinyML by operational constraints and quantitative trade-offs
- Apply the decision framework to select paradigms by privacy, latency, compute, and cost constraints
- Analyze hybrid patterns that combine paradigms to satisfy system constraints
- Evaluate deployment decisions, common fallacies, and universal principles that transfer across scales

2.1 Deployment Paradigm Framework

CONSIDER TWO extremes: a wake-word detector on a smartwatch and a recommendation engine in a data center. The wake-word detector represents a TinyML workload operating under milliwatt power budgets and kilobyte memory limits; the recommendation engine exemplifies a cloud ML workload requiring terabytes of embedding tables and megawatt-scale infrastructure. These systems solve different problems under opposite physical constraints, and their supporting infrastructure differs sharply in scale and design. This reality transforms deployment from an operational afterthought into a first-order engineering decision, one that the D-A-M taxonomy shown in figure 1.4 helps us reason about by foregrounding infrastructure alongside data and algorithms.

Physical constraints determine where an ML model can run and shape what is possible in ways no algorithmic choice can override. Yet deployment is far harder than it appears, and the reason is not the model itself. In production ML systems, the model is often only a small part of the overall system (Sculley et al. 2015). The surrounding infrastructure consists of data collection, feature processing, serving infrastructure, monitoring, and resource management. All of it changes dramatically depending on where the model executes.

The physical constraints that govern each environment (latency, power, and memory) force ML deployment into four distinct paradigms, each with its own engineering trade-offs and system design patterns. **Cloud ML** aggregates computational resources in data centers, offering elastic compute and large-scale storage at the cost of network latency. **Edge ML** moves computation closer to where data originates, reducing remote-path latency and limiting sensitive-data transmission (Shi et al. 2016). **Mobile ML** brings intelligence directly to smartphones and tablets, balancing computational capability against battery life and thermal constraints. **TinyML** pushes intelligence to microcontrollers costing dollars and consuming milliwatts (Janapa Reddi, Plancher, et al. 2022). The categorical envelope from megawatts to milliwatts spans nine orders of magnitude in power, while the representative systems in figure 2.1 span nearly eight. The figure positions five deployment tiers—from ultra-low-power sensors on the left to warehouse data centers on the right—delineating where TinyML, edge AI, and cloud AI operating envelopes overlap.




Figure 2.1: Distributed Intelligence Spectrum: Deployment tiers form a continuum governed by physical constraints, spanning nearly eight orders of magnitude in power. Moving from centralized cloud infrastructure to on-device TinyML trades raw computational capacity and elastic scale for deterministic latency, data locality, and network independence.

Cloud 3 MW
Edge 200 W
Mobile 5 W
Tiny 50 mW
log scale

Power spans the paradigms from megawatts (cloud) to milliwatts (TinyML).

While figure 2.1 illustrates the qualitative topology of distributed intelligence, table 2.1 quantifies the operational envelopes across processing location, round-trip latency, power draw, and memory capacity. This nine-order-of-magnitude power span reflects three inescapable physical boundaries: the speed of light (establishing latency floors), thermodynamic dissipation limits (capping computation per watt), and memory signaling energy (the memory wall).

Table 2.1: The Deployment Spectrum (Conceptual): Four paradigms span nine orders of magnitude in power (MW to mW) and memory (TB to KB). This conceptual overview defines each paradigm by its operating regime; the concrete hardware spectrum later grounds these categories in specific platforms and quantitative decision thresholds. The hardware specifications and physical constants underpinning these numbers are catalogued in the System Assumptions appendix.

Paradigm	Where	Latency (ms)	Power	Memory	Best For
Cloud ML	Data centers	100-500	MW	TB	Training, complex inference
Edge ML	Local servers	10-100	100 W	GB	Real-time inference, privacy
Mobile ML	Smartphones	5-50	3-5 W	GB	Personal AI, offline
TinyML	Microcontrollers	1-10	mW	KB	Always-on sensing

2.1.1 The architectural anchor: The single-node stack

To navigate these operating regimes, we anchor our engineering decisions in a four-layer model of the **single-node stack**, which refines the machine side of the D·A·M taxonomy for one host while data and algorithm enter through workload requirements. At the top of the stack, the application defines the mission through a training-throughput or inference-latency target. Chapter 8 and chapter 13 develop the mechanisms that determine whether the system can meet it. The ML framework translates model code into executable operations (chapter 7). The operating system orchestrates resources and moves data between host memory and accelerator memory. The hardware itself, defined by high-bandwidth memory (HBM) capacity, memory bandwidth, intra-node interconnects such as NVLink, and compute throughput, sets the physical limits, with the memory wall as a recurring constraint (chapter 11).

This stack establishes the chapter's **silicon contract**: the fixed performance bargain a given piece of silicon offers a model, set by memory bandwidth, peak compute rate, and fixed overhead. Every chapter in the first half of this text interrogates one or more of these layers, because understanding how they interact within a single machine is the technical prerequisite for mastering larger distributed scales.

These physical constraints interact with the iron law of ML systems (section 1.7), which decomposes end-to-end latency into data movement, computation, and overhead. Different deployment environments emphasize different constraints: cloud training can be compute bound, mobile systems face power walls, and TinyML devices often reach memory-capacity limits. By pairing the physical constraints with the iron law, we develop a quantitative vocabulary for reasoning about which paradigm fits a given workload and *why*. To anchor this analysis concretely, the chapter introduces five lighthouse models (ResNet-50, GPT-2, DLRM for embedding-heavy recommendation, MobileNetV2, and a keyword spotter for wake-word detection) that span the deployment spectrum and isolate distinct system bottlenecks. These reference workloads recur throughout the book, providing a consistent basis for comparing optimization techniques across chapters, workloads, and deployment tiers.

The physics that creates these paradigm boundaries comes first, followed by the analytical tools (iron law, bottleneck principle, workload archetypes) for mapping workloads to deployment targets. Each paradigm then receives an in-depth treatment covering its infrastructure, trade-offs, and representative workloads, with an eye on how to choose among them when a system could plausibly run in more than one. The chapter closes with a comparative decision framework and the hybrid architectures that combine paradigms when no single deployment target satisfies all requirements.

2.2 Physical Constraints: Why Paradigms Exist

A safety system with a 10 ms reaction budget cannot wait for a cross-country round trip, and a billion-parameter model cannot be squeezed into a microcontroller by better code alone. These

are not implementation bugs; they are consequences of the physical laws of speed of light, power thermodynamics, and memory signaling. Where a system runs reshapes the silicon contract between model and hardware. Three constraints govern the engineering trade-offs ahead: the light barrier, the power wall, and the memory wall.¹

1 | Deployment paradigm:

A distinct operating regime whose boundaries are set by physics, not convention. The Cloud-to-TinyML spectrum spans nine orders of magnitude in power because thermodynamic and electromagnetic constraints create hard walls that no software optimization can cross, forcing qualitatively different system architectures at each tier. Misidentifying the paradigm boundary wastes engineering effort: optimizing a cloud model for 5 percent higher throughput is pointless if the application's 10 ms latency budget demands edge deployment.

2 | Latency:

The time between issuing a request and receiving a result, corresponding to T in the iron law. The light barrier makes this floor irreducible: the speed of light in fiber imposes a ~36 ms minimum round trip across the continental US, consuming the entire latency budget of a 10 ms safety-critical system before any computation begins. Every millisecond consumed by distance is a millisecond unavailable for model inference, which is why the light barrier forces paradigm selection rather than mere optimization.

3 | Dennard scaling: Named after Dennard et al. (1974) at IBM, who described MOSFET scaling relationships under which shrinking devices could reduce voltage and current while keeping power density approximately controlled. As voltage scaling slowed and energy became a first-order architectural constraint, performance growth shifted toward parallelism and specialization: multi-core processors, GPUs, and TPUs (Hennessy and Patterson 2019; Esmaeilzadeh et al. 2011).

The light barrier

The **light barrier** establishes the absolute latency² floor. Let Distance denote the one-way path length and c_{fiber} the propagation speed in optical fiber. Equation 2.1 gives the minimum round-trip time:

$$L_{\text{lat,min}} = \frac{2 \times \text{Distance}}{c_{\text{fiber}}} \approx \frac{2 \times \text{Distance}}{200,000 \text{ km/s}} \quad (2.1)$$

For optical fiber, $c_{\text{fiber}} \approx 200,000 \text{ km/s}$, roughly two-thirds of its vacuum value because light propagates more slowly through glass than through a vacuum.

California to Virginia (~3,600 km straight-line) requires ~36 ms round-trip before any computation begins. Actual services add routing, protocol, queuing, and processing latency. Applications requiring sub-10 ms response *cannot* use distant cloud infrastructure—physics forbids it. This constraint creates the need for edge ML and TinyML: when latency budgets are tight, computation must move closer to the data source.

The power wall

The **power wall** emerged because thermodynamics limits how much computation can occur in a given volume. Dynamic CMOS power follows equation 2.2, where C is effective capacitance, V is voltage, and f is clock frequency. Holding effective capacitance C constant over a local operating range where voltage rises proportionally with frequency ($V \propto f$), dynamic power rises as f^3 . This cubic relationship arises when pushing clock frequencies beyond nominal bounds, requiring higher supply voltages to maintain transistor switching speeds. When voltage scaling stalls entirely at nominal limits, dynamic power scales linearly with frequency, but heat dissipation constraints still prevent continuous clock scaling.

$$\text{Power} \propto C \times V^2 \times f \quad \text{where } V \propto f \implies \text{Power} \propto f^3 \quad (2.2)$$

Under this illustrative assumption, doubling clock frequency requires approximately $8\times$ more dynamic power. The breakdown of broad voltage scaling³ ended “free” frequency speedups and forced the industry toward parallelism (multi-core) and specialization (GPUs, Tensor Processing Units (TPUs)). Mobile devices hit hard thermal limits at 3–5 W; exceeding this causes “throttling,” where the device reduces performance to prevent overheating. In practice, this means a mobile model that runs at 60 FPS for 1 minute may throttle to 15 FPS as the device heats up. This physical limit gives rise to mobile ML: battery-powered devices cannot simply run cloud-scale models locally.

The memory wall

The **memory wall** (Wulf and McKee 1995) reflects the widening bandwidth⁴ gap. A simple book-level sketch uses representative annual growth factors to show why the gap compounds:

$$\frac{\text{Compute Growth Rate}}{\text{Memory Bandwidth Growth Rate}} \approx \frac{1.6}{1.2} \approx 1.33 \quad (2.3)$$

In equation 2.3, the numerator and denominator are dimensionless annual growth factors. A ratio above 1 means compute capability is pulling away from memory bandwidth, so each hardware generation makes data movement a larger share of the performance problem unless the workload increases locality or arithmetic intensity.

The representative trends behind this ratio are that processors have doubled in compute capacity roughly every 18 months, but memory bandwidth has improved only ~20 percent annually. This widening gap makes data movement a frequent bottleneck and energy cost in ML workloads. The constraint affects all paradigms but is especially acute for TinyML, where devices have only kilobytes of memory to work with. Section 11.5.1 develops the architectural responses to the memory wall, including HBM and on-chip SRAM hierarchies.

📌 Checkpoint 2.1: Physical constraints and deployment

Deployment choices are governed by physics, not just preference. Check your understanding:

- ❑ **Light barrier:** can you explain why the speed of light makes cloud ML impossible for < 10 ms safety tasks?
- ❑ **Power wall:** can you explain why thermodynamics (heat dissipation) prevents data center models from running on mobile devices?
- ❑ **Memory wall:** can you explain why data movement is often more expensive (in time and energy) than computation?

These physical laws explain *why* the four paradigms exist. Physics creates the boundaries; privacy regulation, economic incentives, and data sovereignty requirements reinforce and sharpen them without creating them. No regulation can make the speed of light faster, and no economic model can repeal thermodynamics.

Knowing that these barriers exist is necessary but not sufficient. Given a specific ML workload (say, a recommendation engine or a wake-word detector), we need to determine which paradigm fits and which barrier the workload will hit first. The answer requires analytical tools that connect workload characteristics to these physical constraints: the iron law to decompose latency, the bottleneck principle to identify the dominant constraint, and a set of workload archetypes to classify where each model falls on the spectrum.

2.3 Analyzing Workloads

Given a recommendation engine or wake-word detector, the workload question is which term will bind first on the target hardware. The central analytical tool for answering that silicon-contract question is the iron law of ML systems, established in section 1.7. Here, T is total time, D_{vol} is bytes moved, BW is bandwidth, O is total operations, R_{peak} is peak compute rate, η_{hw} is hardware efficiency, and L_{lat} is fixed overhead; equation 2.4 adds these costs for a serialized execution path:

$$T = \frac{D_{\text{vol}}}{BW} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}} \quad (2.4)$$


Equation 2.4 decomposes total latency into three terms: data movement (D_{vol}/BW), compute ($O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$), and fixed overhead (L_{lat}). For a single inference, these costs simply add up—each is paid sequentially. In production systems, however, tasks are processed continuously as a stream, and the analysis shifts from single-task latency to identifying which term limits the system. The answer depends entirely on the deployment environment: a model that is **compute bound** during training may become memory bound during inference; a system that runs efficiently in the cloud may hit power limits on mobile devices. To determine which term dominates, we need a companion principle.

2.3.1 The bottleneck principle

The iron law tells us the cost of each term. The **bottleneck principle** tells us which term matters. In a pipelined ML workload, identifying the system bottleneck matters because optimizing fast operations yields zero benefit while the slowest stage remains unchanged. Modern accelerators use **pipelined execution** to overlap data movement with computation: while the accelerator computes on batch n , the memory system prefetches batch $n+1$. When data movement, computation, and I/O are fully overlapped in steady state, let T_{network} denote remote-transfer time and $T_{\text{bottleneck}}$ the resulting per-task time, including fixed overhead. The faster stages hide behind the slowest overlapped stage, so the iron law's sum becomes a maximum, as equation 2.5 formalizes:

$$T_{\text{bottleneck}} = \max \left(\frac{D_{\text{vol}}}{BW}, \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}, T_{\text{network}} \right) + L_{\text{lat}} \quad (2.5)$$

4 | Memory bandwidth (the memory wall): The term “memory wall” was coined by Wulf and McKee in 1995, who predicted that the processor-memory performance gap would eventually dominate system performance—a prediction that proved prescient for ML workloads in which weight or activation movement can bind performance. In the iron law, bandwidth (BW) appears in the denominator of the data term D_{vol}/BW , so doubling the bytes moved without a corresponding bandwidth increase raises the data-movement time. The resulting bottleneck depends on arithmetic intensity, batch size, data reuse, access pattern, and hardware balance rather than model size alone.



Compute capacity outruns memory bandwidth; the widening gap is the memory wall.

- $\frac{D_{\text{vol}}}{\text{BW}}$ (**Memory**): Time to move data between memory and processor.
- $\frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}$ (**Compute**): Time to execute calculations.
- T_{network} : Time for network communication (if offloading).
- L_{lat} (**Overhead**): Fixed latency (kernel launch, runtime overhead).

When a system is memory bound ($D_{\text{vol}}/\text{BW} > O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$), faster processors (R_{peak}) yield no speedup, just as widening a highway cannot help traffic constrained by a narrow bridge. Engineers must identify the dominant term before optimizing. When network transfer is a candidate term, offloading adds energy to the time-based analysis. The local-versus-offload decision must therefore weigh joules as well as milliseconds.

Napkin Math 2.1: The energy of transmission

Problem: Should a battery-powered sensor process data locally (on-device NPU inference) or send it to the cloud when the energy of transmission collides with the battery-driven energy wall?

Given:

- **Data** (D_{vol}): 1 MB (illustrative payload volume).
- **Transmission energy** (E_{tx}): 100 mJ/MB (Wi-Fi/LTE).
- **Compute energy** (E_{op}): 0.1 mJ/inference (illustrative MobileNetV2 neural-processing-unit anchor).

Math:

1. **Cloud approach:** $E_{\text{cloud}} \approx D_{\text{vol}} \times E_{\text{tx}} = 1 \text{ MB} \times 100 \text{ mJ/MB} = 100 \text{ mJ}$.
2. **Local approach:** $E_{\text{local}} \approx \text{Inference} = 0.1 \text{ mJ}$.

Systems insight: Transmitting raw data is 1,000× more expensive than processing it locally. Even if the cloud had infinite speed ($T \approx 0$), the stated battery budget favors local inference over raw-data transmission. The machine constraint (battery) shapes the deployment choice.

The iron law's variables interact differently across deployment scenarios. Before specific workload archetypes are examined, these core performance determinants need a compact definition.


Systems Perspective 2.1: The iron law as deployment diagnostic

The iron law introduced in section 1.7 is the deployment diagnostic used throughout this chapter: it expresses the total time T required for a workload as the sum of data movement, arithmetic, and latency:

$$T = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}}$$

This decomposition is diagnostic: it quantifies how data volume (D_{vol}), compute capacity (R_{peak}), and fixed latency (L_{lat}) jointly set a workload's time budget. Unlike Amdahl's Law, which focuses on parallel speedup, the iron law binds model work to data movement and fixed latency at the deployment boundary. The frequent misconception is that these terms are independent; in reality they are trade-off axes. Increasing batch size can improve duty cycle and data reuse, shifting a memory-bound kernel toward compute-bound behavior, while also increasing the working set and batch service time.

The iron law quantifies the *cost of each ingredient*; the bottleneck principle identifies the *speed of the assembly line*. As a rule of thumb, use the *additive form* in equation 2.4 when analyzing the latency of a single task, and the *max form* in equation 2.5 when analyzing the throughput of a continuous stream of tasks.



The machine axis sets the deployment envelope.

2.3.2 Workload archetypes

The bottleneck principle reduces optimization to identifying the dominant constraint for a given workload. The answer depends on the D·A·M taxonomy in figure 1.4, which decomposes every ML system into Data, Algorithm, and Machine. Different deployment environments create different bottlenecks along these axes, so the same model family can require different engineering when it moves from a cloud server with terabytes of memory to a microcontroller with kilobytes. The iron law turns that diagnostic into four **workload archetypes**.⁵ These are not model categories; they are recurring physical bottlenecks that determine which engineering moves can help.

The first split separates arithmetic-bound systems from data-movement-bound systems. A **compute beast** performs many calculations per byte loaded, so progress comes from higher arithmetic throughput, better utilization, and more parallel execution; large neural-network training is the canonical case. A **bandwidth hog**, by contrast, waits on dense weight or activation movement, so wider memory buses and better data reuse matter more than additional peak FLOP/s; autoregressive text generation illustrates this regime.

The second split covers workloads where the binding constraint is not dense arithmetic at all. **Sparse scatter** workloads are dominated by irregular table lookups and poor cache locality, so memory capacity, access latency, and communication shape performance in recommendation systems with massive embedding tables. **Tiny constraint** workloads face the opposite envelope: energy per inference and memory footprint, not raw speed, determine whether always-on sensing can run at all.



Lighthouse 2.1: Five reference workloads

The five lighthouse models summarized in table 2.2 recur throughout this book as concrete workloads that span the deployment spectrum and isolate distinct system bottlenecks. Chapter 6 provides full architectural details and model biographies.

Table 2.2: Five Lighthouse Models: Recurring workloads used throughout the book to ground the iron law in concrete practice. Each lighthouse pairs an archetype (Compute Beast, Bandwidth Hog, Sparse Scatter, Tiny Constraint) with the deployment paradigm where it predominantly runs, isolating a distinct systems bottleneck.

Lighthouse	Archetype	Deployment Paradigm
ResNet-50	Compute Beast	Cloud training, edge inference
GPT-2/Llama	Bandwidth Hog	Cloud inference
DLRM	Sparse Scatter	Cloud only (distributed)
MobileNetV2	Compute Beast (efficient)	Mobile, edge
Keyword Spotting (KWS)	Tiny Constraint	TinyML, always-on

These archetypes map naturally to deployment paradigms. Compute beasts and sparse scatter workloads gravitate toward cloud ML where resources are abundant, bandwidth hogs span cloud and edge depending on latency requirements, and tiny constraint workloads belong to TinyML. Each archetype therefore maps to a specific model that recurs throughout this book as one of five reference workloads.

The five lighthouse models ground the abstract interdependencies of the iron law in concrete practice. These summaries connect each workload to the specific iron law bottleneck it exemplifies.

The first lighthouse, **ResNet-50**, classifies images into 1,000 categories, processing each image through approximately 8.2 GFLOP using 25.6 million parameters (102.4 MB at FP32) (He et al. 2016a). Used in medical imaging diagnostics, autonomous vehicle perception pipelines, and as the backbone for content moderation systems, its regular, compute-dense structure makes it the canonical benchmark for hardware accelerator performance.

The language models **GPT-2/Llama** exemplify autoregressive text generation (Radford et al. 2019; Touvron, Lavril, et al. 2023). These models generate text one token at a time. During low-batch decoding, producing each token may require moving most or all model parameters (1.5 billion

⁵ | **Workload archetype:** A classification of ML workloads by their dominant iron law bottleneck rather than their model family. The distinction matters because the optimization strategy differs fundamentally: a compute-bound workload benefits from faster arithmetic (R_{peak}), while a bandwidth-bound workload benefits from wider memory buses or fewer bytes moved through reuse, fusion, compression, or lower precision. Misidentifying the archetype wastes optimization effort on the wrong term of the iron law, as when teams add accelerator FLOP/s to a memory-bound inference pipeline and observe zero speedup.

parameters for GPT-2, 7 billion–70 billion parameters for Llama) through memory. This sequential access pattern creates an autoregressive bottleneck that shapes serving cost (Pope et al. 2023).

The recommendation lighthouse, **DLRM**, represents the embedding-heavy recommendation workload behind large-scale “You might also like” systems (Naumov et al. 2019). It maps users and items to embedding vectors stored in tables that can exceed 100 GB of embeddings, making memory capacity rather than computation the binding constraint.

The mobile lighthouse, **MobileNetV2**, is designed for efficient mobile vision tasks such as classification, detection, and segmentation (Sandler et al. 2018). It performs the same image classification task as ResNet but uses depthwise separable convolutions, which separate spatial filtering from channel mixing, to reduce computation by 13.7 $\times$, enabling real-time inference on smartphones at 3–5 W.

The TinyML lighthouse, **keyword spotting (KWS)**, represents the always-on sensing archetype. KWS systems detect short trigger phrases with compact models built for resource-constrained microcontrollers (Y. Zhang et al. 2017); in the Smart Doorbell scenario, that same pattern becomes a local “Ding Dong” or “Hello” trigger. The lighthouse workload has approximately 200K parameters and fits in about 800 KB, placing it in the kilobyte-memory, milliwatt-budget TinyML regime.

Deployment feasibility is evaluated hierarchically, verifying physical memory fit before checking execution latency. In figure 2.2, horizontal bars plot resource demand as a percentage of supply against the vertical 100 percent saturation threshold for the ESP32-S3 microcontroller, comparing Level 1 memory capacity with Level 2 latency under the service level agreement (SLA).




Figure 2.2: Smart Doorbell Constraint Scorecard: Each bar is demand over supply, so the dashed line marks the point where a level is exactly saturated. The model occupies a small fraction of the ESP32-S3’s 4 MB memory budget (Level 1: PASS), but its 101 ms baseline latency is twice the deliberately tightened 50 ms interaction target (Level 2: FAIL), even though it clears the scenario’s canonical 200 ms budget.

The range in compute requirements and memory footprints explains why no single deployment paradigm fits all workloads. A keyword spotter can operate with roughly 20 MFLOP and 800 KB, while ResNet-50 requires about 8.2 GFLOP and roughly 102.4 MB per image. The reference DLRM example already reaches 100 GB, and production DLRM-style recommendation systems can exceed 100 TB. Language models add a bandwidth-dominated regime: billions of parameters streamed repeatedly from memory during autoregressive inference. These five lighthouse models serve as concrete anchors throughout the book, each isolating a distinct system bottleneck revisited in every chapter.

Analytical tools alone remain abstract until grounded in real silicon. The next step translates the iron law, bottleneck principle, and workload archetypes into quantitative engineering decisions by examining how system balance (the interplay of compute, memory, and I/O) varies across real hardware platforms.

2.4 System Balance and Hardware

Physical constraints translate latency-vs-throughput trade-offs into engineering decisions through concrete numbers. Table 2.3 provides order-of-magnitude latencies that should inform every deployment decision—spanning eight orders of magnitude from nanosecond compute operations to hundreds of milliseconds for cross-region network calls. Detailed hardware latencies and bandwidth constraints are covered in chapter 11. An operation with latency $> X$ *cannot* appear on the critical path of a system whose latency budget is X ms.⁶

Table 2.3: Latency Numbers for ML System Design: Order-of-magnitude latencies across compute, memory, network, and ML operations that determine deployment feasibility. Spanning eight orders of magnitude, from nanosecond compute operations to hundreds of milliseconds for cross-region network calls, these physical constraints shape architectural decisions. For a comprehensive quick-reference including energy ratios and scaling rules, see section D.1.

Operation	Latency	Deployment Implication
Compute		
GPU matrix multiply (per op)	~1 ns	Per-operation latency is small
NPU inference (MobileNetV2)	5–20 ms	Mobile can do real-time vision
LLM token generation	20–100 ms	Perceived as “typing speed”
Memory		
L1 cache hit	~1 ns	Keep hot data in registers
Uncached GPU memory read	~200–500 ns	Hide latency; coalesce accesses
DRAM read (mobile)	50–100 ns	Repeated reads can limit workloads
Network		
Same data center	0.5 ms	Microservices feasible
Same region	1–5 ms	Edge servers viable
Cross-region	50–150 ms	Poor fit for tight latency budgets
ML Operations		
Wake-word detection (TinyML)	~1–10 ms	Always-on feasible at < 1 mW
Face detection (mobile)	10–30 ms	Real-time at 30 FPS
GPT-4 first token	200–500 ms	User notices delay
ResNet-50 training step	200–400 ms	Throughput-optimized

The four deployment paradigms gain precision when grounded in concrete hardware. While table 2.1 defined the paradigms conceptually, table 2.6 provides specific devices, processors, and quantitative thresholds for selecting deployment targets.⁷ The same nine-order power span, coupled with data center power usage effectiveness (PUE) metrics,⁸ joins a cost spread from \$millions to \$10 to determine which paradigm serves a given workload economically.

These hardware differences translate directly into performance bottlenecks. To understand which constraint dominates in each paradigm, we apply the bottleneck principle (section 2.3.1).

🔍 Systems Perspective 2.2: System balance across paradigms

The pipelined form of the iron law of ML systems from section 1.7 adds BW_{IO} , the storage or network I/O bandwidth, to the canonical compute and memory terms. Equation 2.6 formalizes the resulting lower bound on total execution time T :

$$T \geq \max \left(\frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}, \frac{D_{\text{vol}}}{BW}, \frac{D_{\text{vol}}}{BW_{IO}} \right) + L_{\text{lat}} \quad (2.6)$$

Here, O represents total operations, R_{peak} is peak compute rate, η_{hw} is hardware utilization efficiency, D_{vol} is data volume, BW is memory bandwidth, BW_{IO} is I/O bandwidth (storage or network), and L_{lat} is fixed overhead. The lower bound identifies which resource (compute, memory, or I/O) limits performance; equality requires ideal overlap with no additional stalls. The dominant term varies by workload and deployment regime. Faster accelerators help compute-bound cloud training; memory-bound large language model (LLM) decode and edge workloads benefit more from reducing bytes moved or increasing bandwidth. Table 2.4 works

⁶ | **Critical path:** The longest sequential chain of dependent operations in a pipeline. The decision rule in the triggering sentence is strict: if a 200 ms cross-region network call appears anywhere on the critical path, a system with a 100 ms total budget is guaranteed to fail regardless of how fast every other stage runs. Model inference need not be the longest stage; preprocessing and postprocessing can dominate, making the critical path longer than model execution alone suggests.


⁷ | **ML hardware cost spectrum:** AI infrastructure spans six orders of magnitude in cost, from \$10 microcontrollers to multi-million-dollar accelerator clusters. This million-fold range means deployment paradigm selection is simultaneously a physics decision and an economics decision. Even within individual-device choices, the same accuracy target may be achievable on a low-cost microcontroller only after substantial model reduction, or on an expensive data-center accelerator with a much larger resource budget, with fundamentally different latency, power, and operational cost profiles.

⁸ | **Power usage effectiveness (PUE):** This metric isolates facility energy overhead, including cooling and power distribution, from the energy delivered to computing equipment. Even a PUE of 1.06 leaves substantial noncompute power at data-center scale. TinyML does not require dedicated facility cooling, although sensing, communication, conversion losses, and sleep power still contribute to its complete system-energy budget. These device-level costs sit outside PUE.

through all four paradigms, splitting cloud into training and LLM inference, and section A.5.1 maps each dominant term to the optimizations that work and the ones that are wasted, turning this diagnosis into an action plan.

Table 2.4: Dominant Bottleneck by Paradigm: Which iron-law term limits performance in each deployment paradigm, the physical reason it dominates, and the resulting optimization focus. The dominant term shifts the optimization strategy entirely: cloud training maximizes compute utilization, while LLM inference must attack memory bandwidth.

Paradigm	Dominant Constraint	Why	Optimization Focus
Cloud Training	$O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$ (Compute)	Abundant memory/network; FLOP/s limit throughput	Maximize accelerator utilization, batch size
Cloud LLM Inference	D_{vol}/BW (memory bandwidth)	Sequential generation repeatedly moves model state	Increase reuse; reduce bytes moved
Edge Inference	D_{vol}/BW (memory bandwidth)	Limited local bandwidth; models often memory-bound	Smaller models; fewer memory transfers
Mobile	Energy (implicit)	Battery = $\int \text{Power} \cdot dt$; thermal throttling	Lower precision; duty cycling
TinyML	Model footprint must fit on-chip (capacity)	Kilobyte-scale memory; model must fit on-chip	Tiny models and fixed-point arithmetic



Roofline analysis classifies bottlenecks by comparing a workload’s **arithmetic intensity**—the ratio of floating-point operations performed per byte of memory moved (FLOPs/byte)—against the hardware machine balance point ($R_{\text{peak}}/\text{BW}$) (Williams et al. 2009). The framing is informal here; section D.2.1 derives the model in full, defining arithmetic intensity formally and deriving the ridge point that separates the memory-bound and compute-bound regimes. In that framing, the same ResNet-50 model can shift from compute-bound training behavior at high batch sizes to more memory-sensitive single-image inference at batch=1. Deployment paradigm selection must account for this shift.

This shift between training and inference is critical to understand. Recall the D-A-M taxonomy from figure 1.4: every ML system comprises Data, Algorithm, and Machine. Table 2.5 shows how each component behaves differently depending on whether the system is training (learning patterns) or serving (applying them).

Table 2.5: D-A-M × Phase: The same model imposes starkly different demands on Data, Algorithm, and Machine depending on whether the system is training or serving. When bottlenecks shift unexpectedly, check which phase is currently being optimized.

Component	Training (Mutable)	Inference (Immutable)
Data	Massive throughput: large batches, shuffling, augmentation	Low latency: single samples, freshness, speed
Algorithm	Learning phase: update model parameters from examples	Prediction phase: apply fixed weights to new inputs
Machine	Throughput-optimized: high-bandwidth clusters, large memory	Latency-optimized: edge devices, inference accelerators

A quantitative comparison applies this analysis to ResNet-50 inference on a high-end data center accelerator and a mobile NPU. Treat these as point-in-time hardware anchors: the arithmetic is about compute rate, memory bandwidth, and batch size, while chapter 11 explains the accelerator architecture behind the numbers.

Napkin Math 2.2: ResNet-50 on cloud vs. mobile

Problem: Is ResNet-50 inference compute bound or memory bound on (a) a high-end data center accelerator and (b) a flagship mobile NPU?

Given (from lighthouse models):

- ResNet-50: 8.2 GFLOP per inference, 25.6 million parameters (102.4 MB FP32, 51.2 MB FP16)

Analysis:**(a) Cloud data-center accelerator (batch=1, FP16)**

- Peak compute: 312 TFLOP/s (FP16)
- Memory bandwidth: 2.04 TB/s
- Compute time: $T_{\text{comp}} = \frac{8.20 \times 10^9}{3.12 \times 10^{14}} = 0.026 \text{ ms}$
- Memory time (weights plus activations): $T_{\text{mem}} = \frac{5.63 \times 10^7}{2.04 \times 10^{12}} = 0.028 \text{ ms}$
- **Result:** Bottleneck = Memory (1.1× slower than compute)
- **Analysis:** Compute per byte moved = $\frac{8.20 \times 10^9}{5.63 \times 10^7} = 146 \text{ FLOP/byte}$. This ratio measures how much arithmetic the workload performs for each byte loaded. When the ratio exceeds the hardware's compute-to-bandwidth ratio ($R_{\text{peak}}/\text{BW}$), the workload is compute bound; below it, the workload is memory bound. For single-image inference, the low batch size yields limited reuse, explaining why even powerful accelerators can be memory bound at batch = 1.

(b) Mobile: Flagship NPU (batch=1, INT8)

- Peak compute: ~35 TOPS (INT8)—representative of modern mobile NPUs
- Memory bandwidth: ~51.2 GB/s (LPDDR5)
- Model size: 25.6 MB (INT8 quantized)
- Compute time: $T_{\text{comp}} = \frac{8.20 \times 10^9 \text{ INT8 ops}}{3.50 \times 10^{13} \text{ INT8 ops/s}} = 0.23 \text{ ms}$
- Memory time (weights plus activations): $T_{\text{mem}} = \frac{2.82 \times 10^7}{5.12 \times 10^{10}} = 0.55 \text{ ms}$
- **Result:** Bottleneck = Memory (2.3× slower than compute)

Systems insight: Both platforms are memory bound for single-image inference. The A100's faster memory bandwidth (2.04 TB/s vs. 51.2 GB/s = 39.8×) translates to roughly 19.9× faster inference; peak compute alone is not the limiting comparison. This explains why byte-reduction and lower-precision techniques can beat simply buying more peak FLOP/s for deployment. ResNet-50 becomes compute bound when batching and data reuse raise computation per byte moved above the hardware balance point, $I > R_{\text{peak}}/\text{BW}$, where $I = O/D_{\text{vol}}$. The crossover is architecture- and implementation-dependent because activation traffic, input/output movement, cache reuse, and runtime execution details all change the effective bytes moved per inference.

Compression matters more, not less, as the hardware budget tightens. As systems transition from Cloud to Edge to TinyML, available resources decrease dramatically. Table 2.6 quantifies this progression with concrete hardware examples: memory drops from 128 TB (cloud) to 512 KB (TinyML), a 250 million-fold reduction, alongside the same megawatt-to-milliwatt power span. This resource disparity is most acute on microcontrollers, the primary hardware platform for TinyML, where memory and storage capacities are insufficient for conventional ML models.

Table 2.6: Hardware Spectrum (Concrete Platforms): Representative devices that instantiate each deployment paradigm from table 2.1. Where the conceptual table defines operating regimes, this table provides the specific processors, memory capacities, power envelopes, and price points that practitioners use to match workloads to hardware. The DGX Spark sits at the high end of the edge spectrum; most edge deployments use far smaller devices (for example, Jetson Orin Nano). We include it to illustrate the ceiling of noncloud deployment. NVIDIA's Jetson family itself spans a wide SKU spectrum, from Jetson Orin Nano (7–15 W) through Jetson Orin NX (10–25 W) to Jetson AGX Orin (15–60 W); throughout this book, Jetson power figures should be read against the specific SKU named in context.

Category	Example Device	Processor	Memory	Storage	Power	Price Range
Cloud ML	Google TPU v4 Pod	4,096 TPU v4 chips, 1.1 EFLOP/s BF16	128 TB HBM2	Cloud-scale (PB)	~3 MW	Cloud service (rental)
Edge ML	NVIDIA DGX Spark	GB10 Grace Blackwell, 1 PFLOP/s sparse FP4	128 GB LPDDR5x	4 TB NVMe	~200 W	~\$3,000–\$5,000

Category	Example Device	Processor	Memory	Storage	Power	Price Range
Mobile ML	Flagship Smartphone	Mobile SoC (CPU + GPU + NPU)	8–16 GB	128 GB-1 TB	3–5 W	\$999+
TinyML	ESP32-S3	Dual-core @ 240 MHz	512 KB SRAM	4 MB Flash	0.05 W–1.2 W active board power	\$10

The hardware spectrum turns that bottleneck diagnosis into a fit test. The platforms in table 2.6 are products of decades of hardware evolution, from floating-point coprocessors in the 1980s through graphics processors in the 2000s to today’s domain-specific AI accelerators. Chapter 11 traces this historical progression and the architectural principles that drove it. Here, the consequence matters most: qualitatively different hardware appears at different points in the infrastructure, so each workload must be matched to the region whose compute, memory, power, and cost envelope it can satisfy.

Each paradigm occupies a distinct region of the deployment spectrum, governed by the physical constraints (light barrier, power wall, memory wall) and quantified by the analytical tools (iron law, bottleneck principle) introduced in section 2.3. The quantitative thresholds in table 2.7 help practitioners determine whether a workload fits a target’s compute, bandwidth, power, and latency envelope.

Table 2.7: Deployment Decision Thresholds: Practical envelopes that practitioners use to determine deployment feasibility for each paradigm in table 2.6. These values answer the question “can my workload run here?” by specifying the compute tier, memory bandwidth, and power envelope that each paradigm provides. Each power figure marks a paradigm’s typical class rather than a hard ceiling; the edge envelope in particular extends from low-power embedded modules up to workstation-class devices.

Paradigm	Compute	Memory bandwidth	Power	Latency (ms)
Cloud ML	> 1000 TFLOP/s	> 1000 GB/s	MW class (PUE 1.1–1.3)	100-500
Edge ML	~1 PFLOP/s	> 270 GB/s	100 W class	10-100
Mobile ML	15–45 TOPS	51.2–77 GB/s	3–5 W	5-50
TinyML	< 1 TOPS	—	< 1 mW always-on average target	1-10

Table 2.7 completes the fit test: each paradigm in the table is an operating envelope with a different binding resource. The next four sections progress from cloud to TinyML, tracing the gradient from maximum computational resources to maximum efficiency constraints while keeping the same questions in view: which term binds, what optimization helps, and which trade-offs follow.

2.5 Cloud ML: Computational Power

Consider what it took to train GPT-3:⁹ 3,634.3 PFLOP-days of computation, 10,000 V100 GPUs running for approximately 15 days, consuming megawatts of power (Patterson et al. 2021). No smartphone, edge server, or single machine could have completed this training run on a practical schedule. It required a data-center-scale cluster with enough aggregate compute, memory, and storage. This is the defining proposition of cloud ML: when latency can be tolerated, it offers computational scale that no other paradigm can match.

Cloud ML aggregates computational resources in data centers¹⁰ to handle computationally intensive tasks: large-scale data processing, collaborative model development, and advanced analytics. This infrastructure serves as the natural home for three of the four workload archetypes: Compute Beast workloads like ResNet training that demand sustained TFLOP/s across thousands of accelerators, Bandwidth Hog workloads like large language model inference that benefit from TB/s HBM bandwidth, and Sparse Scatter workloads like recommendation systems that require terabytes of embedding tables and high-bandwidth interconnects for all-to-all communication patterns.

Cloud deployments range from single-machine instances (workstations, multi-GPU servers, DGX systems) to large-scale distributed systems spanning multiple data centers. This book focuses

⁹ | **Large language model (LLM) training scale:** GPT-3 required approximately 3,634.3 PFLOP-days (Patterson et al. 2021). This scale illustrates the core cloud ML trade-off: only centralized infrastructure can aggregate enough R_{peak} for large training runs, but remote inference adds network latency and is unsuitable when the end-to-end path exceeds the application’s latency budget.

¹⁰ | **Cloud as utility computing:** The utility model allows providers to offer a specialized hardware portfolio that is economically infeasible for a single organization to maintain. This provides direct, on-demand access to the specific architectures required by each workload archetype: dense accelerator pods for Compute Beasts, HBM-equipped nodes for Bandwidth Hogs, and high-memory systems with fast interconnects for Sparse Scatter. A team can therefore rent a purpose-built, \$10M+ supercomputing pod for a few hours rather than owning it.

on single-machine cloud systems, where the reader learns to build and optimize ML systems on individual powerful machines. Future studies can address distributed cloud infrastructure, where systems coordinate computation across multiple networked machines. This follows the principle of establishing foundations before adding complexity.

Cloud deployment offers elastic compute while adding distance to remote paths.

Definition 2.1: Cloud ML

Cloud machine learning is the deployment paradigm that trades latency for elastic compute by locating ML workloads in centralized data centers, decoupling computational capacity from the physical location of data sources and users.

1. **Significance:** Cloud deployment dominates the R_{peak} term: a single cloud region can provision thousands of accelerators on demand, delivering aggregate throughput that no typical on-premise installation can match economically. The trade-off is the L_{lat} term: a minimum round-trip latency of 10–100 ms (set by the speed of light over continental distances) makes cloud infeasible for any workload requiring sub-10 ms response.
2. **Distinction:** Unlike edge ML, which prioritizes latency determinism and data locality at fixed R_{peak} , cloud ML prioritizes elastic R_{peak} at the cost of variable L_{lat} .
3. **Common pitfall:** A frequent misconception is that cloud ML is “unlimited compute.” In reality, the distance penalty (L_{lat}) and the ingestion bottleneck (D_{vol}/BW) are physics constraints that no software optimization can eliminate, setting a hard floor on response time for any workload whose data originates outside the data center.

Centralization defines the cloud bargain: pooling elastic compute across warehouse-scale clusters eliminates local capacity limits, but introduces distance and network latency. The four-column tree in figure 2.3 maps this trade-off from left to right, connecting the binding scale constraint to its system response, failure boundaries, and well-matched workloads. The examples that thrive in this regime, including virtual assistants, recommendation systems, and fraud detection, all tolerate that bargain because scale matters more than immediacy. The most fundamental challenge, network latency, is not an engineering limitation but a physics constraint. A quick calculation of the distance penalty makes this concrete.

Napkin Math 2.3: The distance penalty

Problem: Consider a real-time safety monitor for a robotic arm. The safety logic requires a 10 ms end-to-end response time to prevent injury, so the distance penalty imposed by the light barrier matters. The model runs in a high-performance cloud data center 1,500 km away. Can the safety budget be met?

Physics:

1. **Light in fiber:** $\sim 200,000$ km/s.
2. **Round-trip propagation:** $(1,500 \text{ km} \times 2) / 200,000 \text{ km/s} = 15 \text{ ms}$.
3. **Result:** Round-trip propagation alone requires 15 ms, exceeding the 10 ms end-to-end budget (~ 5 ms headroom) before the model performs any inference.

Systems insight: Physics has made cloud ML *impossible* for this application. The model must move to the Edge.

11 | Tensor processing unit (TPU): An application-specific integrated circuit (ASIC) that delivers PFLOP/s-scale throughput by hard-wiring its architecture for the matrix multiplication operations that dominate ML workloads. This extreme specialization trades general-purpose flexibility for a $> 10\times$ improvement in performance-per-watt compared to a general-purpose accelerator on the same ML task. The high cost of deploying these accelerators at data center scale is therefore only economical for massive, sustained ML computation.

2.5.1 Cloud infrastructure and scale

Cloud ML aggregates computational resources in data centers at unprecedented scale. Figure 2.4 captures the physical scale behind this abstraction: a cloud TPU¹¹ data center floor. TPU supercomputer designs organize thousands of specialized accelerator chips into data-center-scale systems that deliver PFLOP/s-to-EFLOP/s reduced-precision throughput (Jouppi et al. 2023). Table 2.6




Figure 2.3: Cloud ML Constraint Map: Centralized infrastructure solves the scale problem for compute beasts, bandwidth hogs, and sparse scatter workloads, but it introduces the distance, dependency, privacy, and cost constraints that determine when cloud ML stops being the right deployment target.

quantifies how cloud systems provide orders-of-magnitude more compute and memory bandwidth than mobile devices, at correspondingly higher power and operational cost. These facilities enable workloads that are impractical on resource-constrained devices, but their remote location introduces critical trade-offs, examined next: network round-trip latency rules out real-time applications, and operational costs scale linearly with usage.

The physical reality of PFLOP/s-scale compute is visible in the infrastructure itself: a single facility floor houses thousands of accelerator chips organized into rows of liquid-cooled racks, each rack consuming kilowatts of power to sustain the aggregate throughput that no individual device can approach.




Figure 2.4: Cloud Data Center Scale: Dedicated accelerator facilities aggregate thousands of specialized TPU chips in liquid-cooled racks to deliver exaFLOP-scale matrix throughput. This centralized infrastructure makes large-scale model training and high-throughput serving viable, but incurs megawatt-scale power consumption and speed-of-light network distance penalties. Image source: (Pichai 2023).

Cloud ML excels at processing massive data volumes through parallelized architectures, enabling training on datasets requiring hundreds of terabytes of storage and PFLOPs of computation, resources that remain impractical on constrained devices. The training techniques covered in chapter 8 and the hardware analysis in chapter 11 explain how practitioners achieve this scale.

The same centralization changes how models are shared and operated. Cloud APIs make trained models accessible worldwide across mobile, web, and IoT platforms. Shared infrastructure enables teams to collaborate with integrated version control, while pay-as-you-go pricing models¹² replace much infrastructure purchasing with usage-based operating expense and scale with demand.

A common misconception holds that cloud ML's vast computational resources make it universally superior. Exceptional computational power and storage do not automatically translate to optimal solutions for all applications. The data gravity invariant in section B.1.1 explains why: as data volume scales, the cost of moving it to compute ($C_{\text{move}}(D_{\text{vol}}) \gg C_{\text{move}}(\text{Compute})$) eventually dominates. The trade-offs listed in figure 2.3 become concrete when we consider where edge and embedded deployments excel: real-time response with sub-10 ms decision-making in autonomous control loops, strict data privacy for medical devices processing patient data, predictable costs through one-time hardware investment vs. recurring cloud fees, or operation in disconnected environments such as industrial equipment in remote locations. The optimal deployment paradigm depends on specific application requirements rather than raw computational capability.

2.5.2 Cloud ML trade-offs and constraints

Cloud ML's advantages carry inherent trade-offs that shape deployment decisions. Latency is the most consequential: remote network latency makes cloud processing unsuitable when the end-to-end path exceeds an application's response budget. Unpredictable response times further complicate performance monitoring and debugging across geographically distributed infrastructure.

Privacy and security pose serious challenges for cloud deployment. Transmitting sensitive data to remote data centers expands the disclosure surface and complicates regulatory compliance. Organizations handling data subject to regulations such as the General Data Protection Regulation (GDPR)¹³ or the Health Insurance Portability and Accountability Act (HIPAA)¹⁴ must select, document, and maintain safeguards appropriate to the data, system, and applicable legal duties. Privacy-preserving approaches can reduce how much sensitive data must leave its original environment, but they complement rather than replace these controls.

Cost management introduces operational complexity requiring TCO¹⁵ analysis rather than naive unit comparisons. A worked cloud vs. edge TCO comparison illustrates the gap between sticker price and true system cost.

For the worked comparison, table 2.8 itemizes the annual GPU, network, load-balancer, and observability costs of an illustrative cloud implementation under public list pricing, and table 2.9 itemizes the corresponding hardware, power, cooling, network, and DevOps labor costs of an on-premise edge-server implementation.

Table 2.8: Cloud Inference Annual TCO: Itemized GPU, network, load-balancer, and observability costs for the fixed-capacity cloud implementation of the ResNet-50-scale vision workload.

Cost Component	Calculation	Annual Cost
GPU inference (A10G)	4 instances $\times$ 8760 h/year $\times$ \$0.75/hr	~\$26,280
Network egress	100 GB/day $\times$ 365 d/year $\times$ \$0.09/GB	~\$3,285
Load balancer	\$0.025/hr + LCU charges	~\$3,723
CloudWatch/logging	Monitoring, alerts	~\$2,000
Total Cloud		~\$35,288/year

Cloud deployments also carry operational constraints. Unpredictable usage spikes complicate budgeting, requiring comprehensive monitoring and cost governance frameworks. Network dependency creates a further constraint: any connectivity disruption directly impacts system availability, particularly where network access is limited or unreliable. Vendor lock-in compounds this problem, as dependencies on specific tools and APIs create portability challenges when transitioning between providers. Organizations must balance these constraints against cloud benefits based on their specific application requirements and risk tolerance. Even with these constraints, cloud ML is well suited to consumer applications that require elastic capacity at global scale.

¹² | **Pay-as-you-go pricing:** A cloud economic model where users pay for accelerator-hours consumed rather than hardware owned. Elastic pricing converts the fixed cost of idle R_{peak} into a variable cost proportional to actual utilization, but sustained 24/7 workloads can cost more on cloud than equivalent on-premises hardware amortized over three years, a crossover that drives the total cost of ownership (TCO) analysis in table 2.8.

¹³ | **GDPR (General Data Protection Regulation):** Article 17 provides a right to erasure when specified grounds apply (European Parliament and Council of the European Union 2016). For an ML system, an erasure request requires tracing where the person's data persists across source stores, derived features, logs, caches, and model artifacts. Retraining may be one remediation, but the regulation does not prescribe re-training every model that encountered the data.

¹⁴ | **HIPAA:** The Security Rule requires regulated entities to implement reasonable and appropriate administrative, physical, and technical safeguards for electronic protected health information, including access and audit controls (U.S. Department of Health and Human Services 2005). The rule is technology-neutral and does not mandate confidential-computing enclaves; concrete controls and their performance overhead depend on the system's risk analysis and implementation.

Table 2.9: Edge Inference Annual TCO: Itemized hardware, power, cooling, network, and DevOps labor costs for the on-premise edge-server implementation, exposing labor as the dominant component of this scenario.

Cost Component	Calculation	Annual Cost
Hardware CapEx	$\$15,000 \div 3 \text{ years life}$	~\$5,000
Power (24/7)	$300 \text{ W} \times 8760 \text{ h/year} \times \$0.12/\text{kWh}$	~\$315.4
Cooling overhead	~30% of power	~\$94.6
Network (fiber)	Fixed line for remote management	~\$1,200
DevOps labor	$0.1 \text{ FTE} \times \$150,000 \text{ salary}$	~\$15,000
Total Edge		~\$21,610/year

 Napkin Math 2.4: Cloud vs. edge TCO

Problem: A vision system serves 1M daily inferences at ResNet-50 scale (10 ms latency, 100 KB response). When all costs are included (GPU hours, network egress, power, cooling, and labor, itemized in table 2.8 and table 2.9), is cloud or on-premises edge deployment cheaper over 3 years?

Math: This scenario provisions a fixed pool of 4 cloud GPU instances and one fixed on-premises server. Most costs are fixed, while cloud egress varies with traffic. The ratio of the two annual totals is therefore not a valid break-even utilization. The valid comparison in equation 2.7 at the stated operating point is the annual cost gap:

$$\text{Annual cost gap} = \text{Cloud annual TCO} - \text{Edge annual TCO} \tag{2.7}$$

At the stated 1M-inference/day operating point, the gap is ~\$35,288/year – ~\$21,610/year = ~\$13,678/year. Multiplying this traffic by the ratio of the annual totals would incorrectly treat provisioned capacity as a per-inference expense.

Result: Edge costs about 38.8 percent less at this operating point. Because the fixed cloud configuration already costs more than the edge total before traffic-dependent egress, these assumptions do not produce a positive traffic crossover. A valid break-even requires demand-scaled cloud resources, an explicit capacity staircase, or a defensible per-query price.

Systems insight: Edge TCO is dominated by labor (~\$15,000 ÷ ~\$21,610/year ≈ 69.4 percent), not hardware. Organizations without existing DevOps capacity should factor in the full cost of maintaining on-premise infrastructure.

2.5.3 Large-scale training and inference

Cloud ML’s computational advantages manifest most visibly in consumer-facing applications that require massive scale. Virtual assistants like Siri and Alexa illustrate the hybrid architectures that characterize modern ML systems: wake-word detection can run on dedicated sub-milliwatt-to-milliwatt hardware directly on the device, enabling always-on listening without continuously activating the main processor; initial speech recognition can run on-device for privacy and responsiveness; and complex natural language understanding and generation can use cloud infrastructure for access to larger models and broader knowledge.

Economics drive this architecture as much as latency. Attempting to process voice interactions for billions of devices entirely in the cloud runs into both an economic and an infrastructure ceiling. Quantifying the voice assistant wall shows both limits at once.

 Napkin Math 2.5: The voice assistant wall

Problem: In this illustrative capacity scenario, 1B voice assistant devices (smartphones, smart speakers, earbuds) each issue 20 queries/day as wake-word traffic. What would the infrastructure scaling cost if all queries were served from cloud ML, and how many dedicated data-center equivalents would peak load require?

15 | **Total cost of ownership (TCO):** Quantifies the gap between sticker price and true system cost by including all direct and indirect costs (power, cooling, labor) over a system’s lifetime, trading upfront capital expense (CapEx) against recurring operational expense (OpEx). For an on-premise GPU, the purchase price is only one component of the three-year TCO, alongside operating costs.

Economic wall: First, the cost of serving wake-word traffic from the cloud.

- **Cloud cost:** ~\$0.50/device/year → 1B devices = \$500,000,000/year. Economically prohibitive under these assumptions for a free feature.
- **TinyML alternative:** 0.1–1 mW local wake-word detection, <\$0.01/device/year. Viable at the modeled scale.

Infrastructure wall: Second, the number of data centers peak load would require.

The economic argument is compelling, and the capacity calculation is instructive:

1. **Query volume:** 1 billion devices × 20 queries/day = 20B queries/day.
2. **GPU demand:** Each query requires ~200 ms of GPU time. Total: 1,111,111 GPU-hours/day.
3. **Scenario capacity:** A data-center equivalent provisioned with 10,000 GPUs provides 240,000 GPU-hours/day.
4. **Average requirement:** ~4.6 dedicated data centers just for voice inference.
5. **Peak reality:** Queries cluster in waking hours (~4.5× peak-to-average), requiring ~20.8 data centers at peak.

Bandwidth upper bound: Third, consider the separate counterfactual of moving the audio itself. If every device streamed audio continuously (16 kHz, 16-bit), each would transmit ~32 KB/s. Across 1 billion devices, the aggregate would be 32 TB/s, an illustrative upper bound rather than the query-based request model.

Systems insight: Cloud-only voice processing is not merely expensive; under these assumptions, it is operationally demanding at global scale. Local wake-word detection is a scalable architecture, not a physical necessity.

The voice assistant pipeline illustrates a core systems principle: deployment decisions are constrained by performance requirements, economic realities, and infrastructure physics. The hybrid approach reduces end-to-end latency relative to pure cloud processing while preserving access to the computational capacity needed for complex language understanding.

Large-scale recommendation engines demonstrate the same cloud bargain in a memory-capacity-bound form. These systems process large datasets using collaborative filtering and embedding-heavy architectures such as DLRM¹⁶ to uncover patterns in user preferences. DLRM exemplifies a memory-capacity-bound workload: its embedding tables, representing millions of users and items, can exceed terabytes in size, requiring distributed memory across many servers just to store the model parameters. Cloud computational resources enable continuous updates and refinements as user data grows.

These applications share a common thread: they trade latency for scale, accepting remote-path delay in exchange for computational resources that are difficult to supply locally. Fraud detection systems analyzing millions of transactions, recommendation engines processing terabytes of embedding tables, and language models generating text one token at a time can all depend on this bargain. Yet as the voice assistant wall demonstrated, there exist applications where no amount of cloud compute can compensate for the physics of distance. When latency budgets drop below what the remote path permits, or when data volumes exceed what networks can carry, the computation must move closer to the data source.

2.6 Edge ML: Latency and Privacy

When a remote path exceeds the workload’s latency budget, cloud inference is infeasible. The distance penalty means cross-region requests add network latency before any computation begins. For an autonomous-vehicle braking path or an industrial safety loop, a 100 ms network delay can consume the entire response budget. The logical engineering response is to move the computation closer to the data source.

Edge ML emerged from this constraint, trading elastic cloud capacity for sub-100 ms latency and local data retention. In archetype terms, edge deployment transforms the optimization target: a Bandwidth Hog workload like LLM inference that is memory bound in the cloud can become

¹⁶ | **Deep learning recommendation model (DLRM):** Meta’s 2019 architecture that exemplifies the “Sparse Scatter” archetype. Embedding tables for production recommendation systems can exceed 100 TB, making DLRM constrained by memory capacity and communication BW rather than raw R_{peak} . This inversion of the typical compute-bound assumption forces specialized cluster designs where memory, not arithmetic, is the scarce resource.

dominated by remote-path latency when offloaded from the edge. Edge hardware with sufficient local memory removes that remote-path penalty, shifting attention back to local memory bandwidth and compute. Recall the iron law from equation 2.6: local processing removes the offload term $D_{\text{vol}}/BW_{\text{IO}}$ from the request path, leaving the same memory-vs.-compute trade-off without the remote network penalty.

This paradigm shift is essential for applications where cloud round-trip delays are unacceptable. Autonomous systems requiring split-second decisions and industrial IoT¹⁷ applications demanding real-time response cannot tolerate a remote path that misses the control deadline. Applications subject to data-sovereignty or privacy constraints may also require local processing rather than transmission to a remote data center. Edge devices, including gateways and IoT hubs, occupy a middle ground in the deployment spectrum, maintaining acceptable performance while operating under intermediate resource constraints.

This locality-first trade-off defines the edge paradigm.



Definition 2.2: Edge ML

Edge machine learning is the deployment paradigm optimized for latency determinism and data locality by locating computation physically adjacent to data sources.

1. **Significance:** It circumvents the distance penalty (L_{lat}) of the cloud, trading elastic scale for a fixed local compute capacity (R_{peak}).
2. **Distinction:** Unlike Cloud ML, which prioritizes throughput, edge ML prioritizes determinism and privacy. Unlike TinyML, edge ML may still use workstation-class accelerators such as general-purpose GPUs (GPGPUs).
3. **Common pitfall:** A frequent misconception is that edge ML refers to a specific hardware class. In reality, it is a location paradigm: it spans from IoT gateways to on-premise servers, unified by physical proximity to the data source.

Moving computation to local servers removes long-haul transmission from the request path at the cost of managing distributed infrastructure. In figure 2.5, the constraint tree traces how physical proximity to sensors dictates local filtering and dedicated accelerator hardware while exposing fixed capacity boundaries.




Figure 2.5: Edge ML Constraint Map: Moving computation near the data removes the network-latency term and reduces bandwidth pressure, but it replaces centralized scale with fixed local capacity, distributed security exposure, and operational complexity across many sites.

The benefits of lower bandwidth usage and reduced latency become stark when we examine real-world data rates. The defining characteristic of edge deployment is less about *where* processing occurs than about *how much data* that location must handle. When the data rate exceeds available

¹⁷ | **Industrial IoT (IIoT):** A domain where latency constraints can be set by physical safety rather than user perception. A safety loop with a 5 ms halt budget cannot admit a 100 ms remote round trip, so its sensing and stop decision must remain local. That choice removes the long-haul network term but leaves the edge system bounded by its local compute capacity (R_{peak}).

network capacity, the resulting bandwidth bottleneck forces processing to the edge regardless of other considerations.

Napkin Math 2.6: The bandwidth bottleneck

Problem: Consider a quality control system for a factory floor with 100 cameras running at 30 FPS with 1080p resolution. Does video streaming become the bandwidth bottleneck, and does edge ML reduce the bandwidth enough to process locally?

Physics:

1. **Raw data rate per camera:** $1920 \times 1080 \times 3 \text{ bytes} \times 30 \text{ FPS} \approx 186.6 \text{ MB/s}$.
2. **Total data rate:** $100 \text{ cameras} \times 186.6 \text{ MB/s} = 18.7 \text{ GB/s}$.
3. **Cloud transfer exposure:** Uploading raw camera feeds is primarily a bandwidth, ingest, storage, and processing problem; per-GB cloud egress charges apply when data is transferred back out of the cloud. If the raw stream were later retrieved at \$0.09/GB data-transfer-out pricing, the transfer charge alone would reach \$4.4M/month.
4. **Network reality:** Even a dedicated 10 Gbps line (1.25 GB/s) cannot carry the load—the workload demands 14.9× more bandwidth than exists.

Systems insight: Physics has made cloud streaming *impossible* for this application. Edge processing is not optional—it is mandatory. If an edge server transmits only defect metadata (1 KB per detection at roughly 20 events/s across the floor), bandwidth falls by about 933,120×.

The bandwidth calculation in napkin math 2.6 reveals why edge processing is preferred for that high-volume sensor scenario. For battery-powered edge devices (wireless cameras, drones, wearables), the constraint can be even more severe: under the payload and radio assumptions in napkin math 2.1, radio transmission costs 1,000× more energy than the local inference comparison. This can make cloud offloading impractical, although the result depends on payload size, radio, duty cycle, and local compute. Figure 2.6 broadens the comparison to five illustrative workload-and-deployment energy anchors.




Figure 2.6: Illustrative Energy per Inference Anchors: An eight-order-of-magnitude energy gap separates battery-constrained TinyML keyword spotting ($\sim 10 \mu\text{J}$) from cloud foundation-model queries ($\sim 1 \text{ kJ}$) on a log scale. This physical divergence explains why always-on edge sensing requires specialized low-power architectures rather than continuous wireless streaming to the cloud. (Plotted values represent pedagogical scenario anchors rather than full-system node model outputs).

2.6.1 Edge ML benefits and deployment challenges

Edge ML spans wearables, industrial sensors, and smart home appliances that process data locally¹⁸ without depending on central servers. The illustrative eight-order span in figure 2.6 reflects different workloads and accounting boundaries, not an irreducible constant of deployment tiers.

¹⁸ | **Internet of things (IoT) data wall:** McKinsey estimates that IoT deployments could create trillions of dollars in economic value by 2030, but those deployments depend on continuous sensor streams from devices distributed across homes, factories, farms, vehicles, and infrastructure (McKinsey Global Institute 2021). In many of those settings, the aggregate D_{vol} from raw streams overwhelms the available uplink budget or latency budget for centralized ingestion, making local edge processing an architectural requirement rather than merely a cost optimization.

The same energy boundary becomes a model-size boundary. Because edge devices operate within tight power envelopes, memory capacity limits what can fit while bandwidth limits how quickly the weights and activations can move. Those constraints motivate the optimization techniques covered in chapter 10, which reduce model bytes and operations to fit the hardware budget. The payoff extends beyond compute: processing raw camera feeds locally can avoid terabit-scale uplink requirements because raw data never leaves the device, reducing recurring cloud-transfer, storage, and processing costs.

2.6.2 The data locality invariant

The decision between local edge processing and remote cloud processing is governed by a bandwidth-latency trade-off. Remote execution is infeasible when its complete path violates the service deadline; otherwise, the faster feasible path follows from comparing complete local and remote times.



Definition 2.3: The data locality invariant

The data locality invariant first tests a complete remote path against the service deadline. Let T_{remote} denote complete remote-path time, T_{deadline} the deadline, BW_{network} offload bandwidth, $L_{\text{lat,network}}$ network latency, and $R_{\text{peak,remote}}$ and $\eta_{\text{hw,remote}}$ the remote compute rate and efficiency:

$$T_{\text{remote}} = \frac{D_{\text{vol}}}{BW_{\text{network}}} + L_{\text{lat,network}} + \frac{O}{R_{\text{peak,remote}} \cdot \eta_{\text{hw,remote}}}, \quad \text{remote feasible} \iff T_{\text{remote}} \leq T_{\text{deadline}}$$

The corresponding T_{local} is the complete local-path time. If both paths are feasible, compare T_{local} and T_{remote} ; either path can be rejected by the deadline.

1. **Significance:** The invariant defines a crossover beyond which adding remote compute (R_{peak}) has little effect because transmission dominates T_{remote} . The local path is preferred only after its own complete execution time and resource limits are checked.
2. **Distinction:** Unlike the iron law, which decomposes one execution path, the locality invariant compares two complete paths against each other and against an explicit service deadline.
3. **Common pitfall:** A frequent misconception is that 5G/6G “solves” locality. While these technologies improve BW_{network} , they do not remove propagation and protocol latency, so a remote path can still miss a tight service deadline.

The locality crossover is easiest to see by comparing a single high-rate sensor frame with the round-trip budget for remote processing.

Physics frames the architectural choice; the engineering trade-offs follow from it. The most immediate benefit is latency: avoiding cloud round trips can reduce response time, enabling safety-critical applications when the complete local path meets its real-time deadline. Bandwidth savings compound this advantage; for example, a retail store with 50 cameras can reduce transmission substantially by processing locally and transmitting only metadata, though the exact reduction and cost depend on codecs, traffic, and service pricing. Privacy strengthens in turn, because keeping raw data local reduces transmission exposure without eliminating other privacy obligations. For industrial deployments, operational resilience can be decisive: systems continue functioning during cloud-link outages when local power and networking remain available, a property essential where downtime carries immediate cost.

These benefits carry corresponding limitations that compound as deployments scale. In this chapter’s edge scenario, hardware provides 1–8 GB of memory and 5–50 W of power, well below the cloud-server reference.¹⁹ These constraints make model footprint, activations, runtime state, and sustainable duty cycle first-class limits. Managing distributed networks introduces complexity that scales nonlinearly with deployment size, because coordinating version control and updates across thousands of devices requires sophisticated orchestration systems,²⁰ and hardware heterogeneity across diverse platforms demands different optimization strategies for each target.

¹⁹ | **Edge server constraints:**

In this chapter’s representative envelope, 1–8 GB can hold roughly 1–8 billion INT8 weights before activations, runtime state, and other buffers, so parameter count alone does not define a universal millions-only ceiling. Sustainable compute is further limited by power, cooling, and duty cycle.

²⁰ | **Edge fleet coordination:**

Managing thousands of distributed edge devices introduces failure modes absent from centralized cloud: intermittent connectivity causes model version drift, hardware heterogeneity requires per-target optimization, and physical accessibility makes firmware rollbacks costly. These operational patterns are examined in chapter 14.

Napkin Math 2.7: The locality crossover

Problem: Should a drone's object avoidance system (4K, 60 FPS) offload to the cloud, or does this become a locality crossover?

Given:

- **Data** (D_{vol}): 4K frame $\approx$ 24.9 MB.
- **Bandwidth** (BW_{network}): 100 Mb/s home broadband (up).
- **Remote network and compute response:** 110 ms (round-trip + remote compute).

Math:

1. **Transmission time:** $24.9 \text{ MB} \times 8 \text{ bits/100 Mb/s} = 1,990.7 \text{ ms}$.
2. **Complete remote path:** Transmission + 110 ms remote response = 2100.7 ms, vs. a 16.7 ms frame interval.

Systems insight: The remote path misses the frame interval before control-loop slack is considered, and faster remote compute cannot remove the transmission time. Cloud offload is therefore infeasible under these assumptions; the local path must still be measured to confirm that it meets the deadline.

A realistic retail deployment shows how those constraints turn into throughput, hardware, and fleet-cost requirements. Consider a smart retail chain that deploys person detection across 500 stores, each with 20 cameras/store running at 15 FPS. Table 2.10 cascades the per-store inference rate through YOLOv8-nano's per-frame FLOP count to yield the throughput each store must sustain.

Table 2.10: Edge Inference Sizing Requirements: Per-store throughput target for the smart-retail person-detection scenario.

Metric	Calculation	Result
Inferences per store	20 cameras/store $\times$ 15 FPS	300 inferences/s
Model compute	YOLOv8-nano: 8.7 GFLOP/inference	2610 GFLOP/s
Required throughput	2610 GFLOP/s $\times$ 2 (headroom)	$\sim$ 5.22 TOPS equivalent

Table 2.11 compares throughput, power, unit cost, and fleet cost for three candidate edge accelerators, including embedded GPU accelerators.

Table 2.11: Edge Accelerator Options: Throughput, power, and cost for three candidate edge accelerators at fleet scale.

Edge Device	INT8 TOPS	Power	Unit Cost	Fleet Cost
NVIDIA Jetson Orin NX	100 TOPS	10–25 W	\$600	\$300,000
Intel NUC + Movidius	1 TOPS	15 W	\$400	\$200,000
Google Coral Dev (3 boards/store)	peak 12 TOPS derated 6 TOPS	6 W	\$450	\$225,000

Napkin Math 2.8: Edge inference sizing

Problem: Given the throughput target in table 2.10 and the candidates in table 2.11, which edge accelerator (USB-scale TPU, workstation-class embedded GPU, or general-purpose mini-PC) delivers the required throughput at the lowest three-year fleet cost?

Math: We size from what one board sustains, not from what its datasheet peaks at. A Coral Dev Board is rated at 4 TOPS peak but delivers about 2 TOPS after 50 percent derating, so the count per store is 5.22 TOPS required $\div$ 2 TOPS per board = 2.6 boards, which rounds up to 3 boards/store.

Over 3 years across 500 stores, the TCO comparison for that configuration is Hardware \$225,000 + Power ($0.006 \text{ kW} \times 500 \times 8760 \text{ h/year} \times 3 \text{ years} \times \$0.12/\text{kWh} = \$9,460.8$) = \$234,460.8 total vs. a one-dedicated-cloud-GPU-per-store baseline at ~\$9,855,000.

Result: One Coral alone is undersized, but the sharded configuration with 3 boards/store is the low-cost edge choice, providing 6 TOPS and about 1.3× lower per-store hardware CapEx than Jetson. Jetson remains the simpler single-device deployment when integration complexity matters more than hardware cost.

Systems insight: Edge sizing is a capacity-and-cost problem, not a device-name problem. The cheapest feasible design may be several small accelerators per site rather than one larger board, but that hardware saving must be weighed against the coordination and maintenance burden of a sharded edge fleet.

Security challenges intensify because edge devices are physically accessible: equipment deployed in retail stores or public infrastructure faces tampering risks that centralized data centers do not, requiring hardware-based protection mechanisms such as secure boot, encrypted storage, and tamper-evident enclosures. In an illustrative deployment with \$500–2,000 of edge hardware per site, instrumenting 1,000 sites requires \$500,000–2,000,000 upfront before installation and operations. Whether that capital cost beats cloud spending depends on utilization, labor, maintenance, and service pricing.

2.6.3 Real-time industrial and IoT systems

Edge applications differ by domain, but each makes the same locality argument concrete: the system cannot wait for the network, cannot ship the data, cannot expose the raw signal, or cannot stop during an outage. Autonomous vehicles represent a demanding application because safety-critical perception and control decisions must fit millisecond-scale budgets. Onboard systems process streams from multiple cameras and other sensors through local hardware; a remote cloud path cannot satisfy the tight control loop when its network delay alone exceeds the deadline.

Smart retail environments demonstrate edge ML's practical advantages for privacy-sensitive, bandwidth-intensive applications. A checkout-free store²¹ can process feeds from many cameras on local edge servers, tracking customer movements and item selections without continuously transmitting raw video over a wide-area link. This architecture reduces bandwidth demand and transmission exposure, but it does not by itself ensure privacy: retention, access, secondary use, and governance controls still determine the system's compliance and risk.

The industrial IoT²² uses edge ML for applications where millisecond-level responsiveness directly impacts production efficiency and worker safety. Manufacturing facilities deploy edge ML systems for real-time quality control and predictive maintenance.²³

Smart buildings use edge ML to optimize energy consumption while maintaining operational continuity during network outages. Commercial buildings equipped with edge-based building management systems process data from thousands of sensors monitoring temperature, occupancy, air quality, and energy usage. This reduces cloud transmission requirements by an order of magnitude or more while enabling sub-second response times. Healthcare applications similarly use edge ML for patient monitoring and surgical assistance, reducing data transmission while supporting low-latency workflows for real-time guidance.

Many industrial and retail edge systems are stationary and connected to site power. Recall the iron law in equation 2.6: local deployment removes the remote network term, while mains power makes the energy budget less restrictive than a battery's. A factory edge server consuming hundreds of watts around the clock may be acceptable when connected to site power. Mobile devices, however, travel with their users and run on fixed battery budgets. When we shift from stationary edge infrastructure to the smartphone in a user's pocket, a new term enters the optimization: $\text{Energy} = \text{Power} \times T$. The dominant constraint can shift from *latency* to *energy per inference*, and with it, the engineering calculus.

21 | Checkout-free retail:

A store may fuse camera observations with shelf events and transaction state locally, then transmit only selected events or aggregates. The privacy boundary must still include those derived records, because reducing raw-video transport does not determine retention, access, or secondary use.

22 | Industry 4.0:

The fourth industrial revolution integrates ML into the sensor-actuator feedback loop on factory floors. The systems consequence is that the control loop latency (L_{lat}) must be shorter than the physical process it governs: a welding robot that detects a defect at 60 Hz has a 16.7 ms sampling interval, requiring a deployment whose complete response path fits that deadline.

23 | Predictive maintenance:

Models analyze sensor data such as vibration and temperature to estimate equipment condition and failure risk. Continuous monitoring makes sensing cadence, duty cycle, power, and local storage first-class constraints. Battery-powered nodes may require sub-watt average operation, whereas mains-powered industrial gateways can use a much larger envelope.

2.7 Mobile ML: Offline Intelligence

Edge ML solves the distance problem that limits cloud deployments by moving computation nearer to the data. Many edge systems are stationary gateways, factory servers, and retail appliances, whereas users carry phones and tablets through changing environments. Bringing the same locality to a portable device introduces a different constraint: the battery. Unlike plugged-in edge servers that can consume hundreds of watts continuously, mobile devices must operate for hours or days on fixed energy budgets.

Mobile ML addresses this challenge by integrating machine learning directly into portable devices such as smartphones and tablets, providing responsive, personalized capabilities. This paradigm excels when data locality, offline operation, and immediate response matter more than computational scale, supporting applications such as voice recognition, computational photography,²⁴ and health monitoring. Local processing reduces transmission exposure but does not by itself guarantee privacy. These battery-powered devices must balance performance with power efficiency and thermal management, making them suited to frequent, short-duration AI tasks.

The mobile environment introduces a critical constraint absent from stationary deployments: energy per inference becomes a first-order design parameter. Under the iron law in equation 2.6, cloud and edge systems optimize for minimizing T , total latency. Mobile systems face an additional constraint: Energy = Power $\times$ T , and the power wall described by equation 2.2 caps sustained power at 3–5 W. In archetype terms, a Compute Beast workload like image classification must be transformed into a more compact vision model,²⁵ reducing FLOPs by 13.7 $\times$ while preserving enough accuracy for the application. This is not merely optimization; it represents a qualitative shift in the compute-per-byte trade-off, accepting lower peak throughput in exchange for sustainable operation within a 3–5 W thermal envelope.

This battery-and-thermal boundary gives the mobile paradigm its defining shape.

Definition 2.4: Mobile ML

Mobile machine learning is the deployment paradigm bounded by thermal design power (TDP) and battery energy.

1. **Significance:** It is constrained by the few-watt heat dissipation capacity of passive cooling, requiring architectures that prioritize sustained energy efficiency over peak throughput (R_{peak}).
2. **Distinction:** Unlike edge ML, which may have active cooling, mobile ML must operate within a personal energy budget. Unlike TinyML, it still provides a rich OS and multi-watt compute capacity.
3. **Common pitfall:** A frequent misconception is that mobile ML performance is a fixed value. In reality, it is a time-varying constraint: performance often drops as the device hits its thermal wall, triggering throttling that reduces the duty cycle (η_{hw}).

On-device processing provides immediate responsiveness and data locality, but operates within strict electrochemical and thermodynamic boundaries. The four-branch taxonomy in figure 2.7 links the phone's shared battery and passive cooling chassis to necessary NPU kernel fusions and thermal throttling limits.

The battery life and resource constraints in mobile ML translate directly into engineering requirements. Always-on ML features incur the **battery tax**, because continuous inference spends the phone's finite energy budget even before the rest of the system runs.

The battery constraint limits total energy consumption over time. Even if we could ignore battery life for a plugged-in tablet or a short demo, thermodynamics imposes a second boundary. Electrical power ultimately becomes heat that the device must dissipate. A data center uses active cooling to remove this heat. A thin mobile device instead relies primarily on passive heat transfer through its chassis and surfaces. This creates the **thermal wall**, a hard ceiling on sustained power consumption that exists independently of battery capacity.

²⁴ | **Computational photography:** Uses ML algorithms (for example, multi-frame fusion, neural denoising) to overcome the physical limits of small mobile camera sensors. This exemplifies the mobile computing trade-off, as a multi-stage pipeline must execute within the user's perceived shutter delay while adhering to a shared 3–5 W thermal budget.

²⁵ | **Mobile vision model reduction:** MobileNet-style architectures reduce the computation in common vision layers while preserving useful accuracy for mobile tasks. The architectural details appear in chapter 6; the systems point here is that mobile deployment often requires changing the model family, not merely moving the same model to a phone.




Figure 2.7: Mobile ML Constraint Map: On-device processing buys responsiveness, privacy, offline operation, and personalization, but the phone's shared battery and passive thermal envelope make sustained energy efficiency the binding design constraint.

Napkin Math 2.9: The battery tax

Problem: Consider deploying a “real-time” background object detector on a smartphone. The model consumes 2 W of continuous power when active. The phone has a standard 15 Wh battery. Can the feature stay on all day? This is the battery tax that turns battery life into a mobile ML energy-budget constraint.

Physics:

1. **Ideal runtime:** $\frac{15Wh}{2W} = 7.5$ hours
2. **Reality:** A user expects their phone to last 24 hours. Running this single feature continuously for a day would require 320 percent of the phone's daily energy budget.

Systems insight: The model cannot simply be “deployed.” The techniques in chapter 10 must reduce both model work and duty cycle so the feature can stay on all day.

The distinction matters for engineering decisions: the battery tax is a budget problem, solvable in principle by reducing how often the model runs or by increasing battery capacity. The thermal wall is a physics ceiling. A larger battery cannot raise the maximum sustained wattage a passive chassis can dissipate; duty-cycle and software changes can only reduce heat generation. Sustained operation above the thermal envelope triggers hardware throttling regardless of remaining battery energy. The two constraints therefore attack different points in the iron law: the battery limits total operations per charge (O integrated over time), while the thermal wall caps the instantaneous rate ($R_{\text{peak}} \cdot \eta_{\text{hw}}$) the silicon can sustain.

Napkin Math 2.10: The thermal wall

Problem: An unoptimized vision model requires 12 W peak compute. Can it be deployed on a mobile device, or does it hit the thermal wall and the mobile power wall?

Scenario: A mobile system-on-chip (SoC) allows approximately 3 W for passive cooling. Assume the unoptimized model ships anyway; a representative throttling trace for such a passively cooled device runs in three steps:

1. **Temperature rise:** At 12 W, the device temperature rises at approximately 1 °C/s.
2. **Thermal trip:** Within 60 s, the hardware reaches the thermal trip point (80 °C), triggering OS throttling.
3. **Throttled throughput:** The 100 FPS model suddenly drops to 30 FPS to stay within the thermal envelope.



Sustained thermal performance falls well below burst peaks.

Math: Assume an illustrative $4\times$ active-power reduction after an FP32-to-INT8 redesign (see chapter 10): $12\text{ W} \div 4 = 3\text{ W}$, which lands right at the 3 W passive-cooling ceiling with no headroom to spare. The factor is a scenario input, not a universal whole-system guarantee from quantization alone.

Systems insight: Even a $4\times$ power reduction does not create thermal headroom; it only just reaches the sustainable limit. Physics sets a hard ceiling that no optimization can exceed.

2.7.1 Mobile ML benefits and resource constraints

Mobile devices exemplify intermediate constraints: 8–16 GB RAM (varying from mid-range to flagship), 128 GB–1 TB storage, 15–45 TOPS AI compute through neural processing units²⁶ consuming 3–5 W power. System-on-chip architectures²⁷ integrate computation and memory to minimize energy costs. Memory capacity constrains model footprint, while 51.2–77 GB/s of bandwidth constrains how quickly weights and activations move. Battery constraints (15 Wh–22 Wh capacity) make energy optimization critical: adding 1 W of continuous ML processing to a phone that otherwise lasts 24 hours would reduce runtime to roughly 9.2 hours–11.5 hours, depending on battery capacity. Specialized on-device ML frameworks provide hardware-optimized inference for 5–50 ms UI response times.

Mobile ML excels at delivering responsive experiences with stronger data locality. Real-time processing can reach sub-10 ms latency for some tasks, enabling 5–50 ms UI response times in interactive applications. Processing sensitive inputs locally reduces data transmission and central storage, and on-device enclaves such as Apple’s Secure Enclave can further protect biometric data,²⁸ though the strength of privacy guarantees ultimately depends on the complete system design and threat model. Navigation, translation, and media-processing features can also operate without a network when their models and supporting data fit the mobile budget. Personalization rounds out the advantage, because models can exploit on-device signals and user context while keeping raw data local.

These benefits require accepting tight resource constraints. Compared to cloud deployments, mobile applications often operate under much tighter memory, storage, and latency budgets, which constrain model size and batch behavior. Battery life presents visible user impact, and thermal throttling can materially limit sustained performance: peak NPU throughput can substantially exceed what the device sustains under prolonged workloads. Development complexity multiplies across platforms, demanding separate implementations and careful performance tuning, while device heterogeneity requires multiple model variants. App-store review and staged release processes can also slow iteration compared with a service controlled entirely by its operator.

2.7.2 Personal assistant and media processing

Across mobile applications, the central systems problem is that several short pipelines must share the same battery and thermal budget. Computational photography exemplifies the challenge of running multiple ML pipelines within a thermal envelope. Modern flagships process each photo through multiple ML stages: portrait mode²⁹ uses depth estimation and segmentation, night mode aligns multiple frames with ML-based denoising, and HDR merging, super-resolution, and scene optimization run in sequence. The engineering challenge is not any individual model but the pipeline: these models must share a 3–5 W power budget and complete within the user’s perceived shutter delay, requiring careful scheduling across CPU, GPU, and NPU to avoid thermal throttling.

Voice-driven interactions demonstrate mobile ML’s layered architecture. Wake-word detection runs continuously on a dedicated low-power core, while speech recognition and keyboard prediction use higher-power neural processing tiers. Each layer operates at a different power tier, illustrating how mobile ML partitions workloads across heterogeneous processing units within a single SoC.

Health monitoring and augmented reality push mobile ML to its sustained-performance limits. Wearables like Apple Watch process ECG and accelerometer data on-device to reduce data transmission, while AR frameworks demand a 16.7 ms total frame budget at 60 FPS for simultaneous localization, hand tracking, and scene understanding. These applications represent the ceiling of what battery-powered, passively-cooled devices can sustain, and they define the boundary beyond which mobile optimization alone is insufficient.

²⁶ | **Neural processing unit (NPU):** A dedicated hardware block on a mobile system-on-chip whose circuits are designed for low-precision tensor operations. This specialization can improve energy efficiency relative to CPU execution, helping AI workloads fit within mobile power budgets.

²⁷ | **System-on-chip (SoC):** Mobile SoCs integrate CPU, GPU, and NPU cores with caches and memory controllers to reduce data-movement costs. Mobile DRAM is generally off-die, and feasible model size depends on capacity, precision, working memory, bandwidth, and latency.

²⁸ | **Face ID:** Apple performs facial matching within the Secure Enclave and encrypts the stored mathematical representations so that they remain available only to that subsystem and do not leave the device (Apple 2024a, 2024b). Apple reports a random false-match probability below 1 in 1,000,000 for one enrolled appearance, with higher probabilities for some populations and configurations. The example shows how local inference and hardware isolation can reduce transmission exposure while supporting low-latency authentication.

²⁹ | **Portrait Mode pipeline:** This is not a single model but a sequence of real-time models for depth estimation, segmentation, and rendering. The core engineering problem is managing the pipeline’s aggregate latency and power, not any single model’s performance. The entire pipeline must execute within the user’s perceived shutter delay and share the phone’s 3–5 W thermal budget, forcing scheduling trade-offs across the CPU, GPU, and NPU to avoid throttling.

These successes can create a misleading sense of ease. A common pitfall involves attempting to deploy desktop-trained models directly to mobile or edge devices without architecture modifications. Models developed on powerful workstations often fail when deployed to resource-constrained devices. A desktop ResNet-50 pipeline may require gigabytes once activations, batches, preprocessing buffers, and runtime overhead are included, even though the FP32 weights alone are about 102.4 MB. Such a pipeline, requiring 8.2 GFLOP per inference, cannot run unchanged on a representative low-end edge target with 512 MB of RAM and a 1 GFLOP/s processor. Beyond simple resource violations, desktop-optimized models may use operations unsupported by mobile hardware, assume numerical formats unavailable on embedded systems, or require batch processing incompatible with single-sample inference. Successful deployment demands architecture-aware design from the beginning: model families, arithmetic formats, and implementation choices must all match the target device.

Mobile ML demonstrates that useful intelligence can operate within a 3–5 W thermal envelope on battery power. However, smartphones still cost hundreds of dollars, require gigabytes of memory, and demand user attention to recharge daily. These requirements make them unsuitable for a vast class of applications: monitoring soil moisture across a thousand-acre farm, detecting structural stress in bridge cables, or listening for endangered species in a remote forest. These scenarios demand not just lower power but a qualitatively different engineering regime, one where the device costs dollars instead of hundreds, memory is measured in kilobytes instead of gigabytes, and the system runs unattended for months or years. Mobile optimization methods help, but they cannot bridge a 10,000-fold gap in available memory. What is needed is not a scaled-down smartphone but an entirely different class of hardware and algorithms.

2.8 TinyML: Ubiquitous Sensing

Imagine instrumenting every pallet in a warehouse, every cable on a suspension bridge, every beehive in an apiary. To put “eyes and ears” on this many physical objects, tens of thousands to millions, the device must cost dollars, not hundreds of dollars, and measure millimeters, not centimeters. Smartphones are far too expensive and too large; what is needed is ubiquitous sensing at the scale of a postage stamp and the price of a cup of coffee.

TinyML completes the deployment spectrum by pushing intelligence to its physical limits: low-cost, low-power ML on deeply constrained embedded devices (Janapa Reddi, Plancher, et al. 2022). In this book’s operating envelope, devices costing less than \$10 and consuming less than one milliwatt³⁰ of power make ubiquitous³¹ sensing economically practical at massive scale; MLPerf Tiny and MCUNet show how such constraints are evaluated and optimized in practice (C. Banbury et al. 2021; Lin et al. 2020). This is the exclusive domain of the Tiny Constraint archetype, where the optimization objective shifts from maximizing throughput to minimizing energy per inference. Under the duty-cycle assumptions used in this chapter, a keyword spotting model consuming 10 μ J per inference can operate for years on a coin-cell battery, achieving million-fold improvements in energy efficiency by trading model capacity for operational longevity.

30 | The 1 mW Threshold: Ambient energy harvesting can power a device indefinitely only when long-term harvested power exceeds the complete system’s average load. Available power varies across sources and environments—solar cells the size of a thumbnail may produce ~ 10 mW outdoors but ~ 10 μ W indoors, thermoelectric generators on warm pipes ~ 100 μ W, and RF energy from nearby transmitters ~ 10 μ W. This crossover can transform the deployment model from “battery-limited lifetime” to “deploy and forget,” but it depends on the source, environment, and workload.

31 | Ubiquitous computing: Mark Weiser’s ubiquitous-computing vision imagined computation woven into everyday environments until it receded from direct attention (Weiser 1991). TinyML is one modern path toward that vision: when the cost and power of an intelligent sensor become low enough for mass deployment, the optimization objective shifts from performance (throughput) to power (energy per inference), the central trade-off of the Tiny Constraint archetype.




Figure 2.8: TinyML Hardware Scale: Microcontroller development boards integrate sensors, processing cores, and low-power radios onto single substrates operating under strict kilobyte-scale SRAM and milliwatt power bounds. This compact footprint shifts the optimization objective from peak throughput to energy per inference, enabling multi-year deployment on coin-cell batteries. Image source: Figure 6 in (Janapa Reddi, Plancher, et al. 2022), licensed under CC BY 4.0.

Where mobile ML requires sophisticated hardware with gigabytes of memory and multi-core processors, TinyML operates on microcontrollers³² with kilobytes of RAM and single-digit dollar price points (C. Banbury et al. 2021; Lin et al. 2020; Janapa Reddi, Plancher, et al. 2022). This radical constraint forces an entirely different approach to machine learning deployment, prioritizing ultra-low power consumption and minimal cost over computational sophistication. TinyML systems power applications such as predictive maintenance, environmental monitoring, and simple gesture recognition. The energy gap between TinyML and cloud inference (figure 2.6) spans at least six orders of magnitude³³ and reaches eight orders for cloud LLM queries, driving entirely different system architectures and deployment models. This extraordinary efficiency enables operation for months or years on coin-cell batteries.³⁴ In figure 2.8, annotated development boards illustrate this physical scale, showing integrated microcontroller chips, sensor arrays, and pin headers designed to run compact models in disconnected environments.

Fitting the entire inference pipeline—from sensor acquisition to classification output—into these diminutive physical packages requires rethinking how model weights and runtime activations are stored in silicon.

Definition 2.5: TinyML

TinyML is the machine learning domain of always-on sensing constrained by kilobyte-scale memory and milliwatt-scale power.

1. **Significance:** It necessitates models whose weights and code fit in flash while activations and runtime state fit in on-chip SRAM, enabling continuous inference on milliwatt power budgets.
2. **Distinction:** Unlike mobile ML, which uses multi-watt processors and a full OS, TinyML runs on microcontrollers using bare-metal firmware or a real-time operating system.
3. **Common pitfall:** A frequent misconception is that TinyML is just “small models.” In reality, it is an energy-bound paradigm: the primary metric is energy per inference (microjoules), not just the parameter count.

TinyML’s milliwatt-scale power consumption creates an energy gap of several orders of magnitude relative to cloud inference, with profound implications for system design. In terms of equation 2.6, TinyML often reaches a memory-fit constraint that the latency equation does not explicitly capture: the model footprint M_{model} and activation footprint must stay within memory capacity C_{mem} . When total memory is measured in kilobytes, model structure and runtime allocation determine feasibility before latency optimization begins, and every avoidable data movement spends part of the energy budget. The objective therefore shifts from raw throughput toward energy per inference and memory fit.

Systems Perspective 2.3: Reading the energy-per-inference numbers

The energy values in table 2.12 combine solver-derived device-energy estimates with separate scale anchors; they are not controlled full-system measurements. An illustrative A100-only ResNet-50 bound is under 1 ms and about ~0.3 J, while an entire server can draw about ~1 kW when CPUs, memory, and overhead are included. The final-column counts show how many inferences one battery’s nominal energy could supply; they are not elapsed battery life and omit sensing, sleep, and communication.

This TinyML energy-efficiency gap helps explain why low-energy local models are attractive for always-on sensing. Actual runtime still depends on inference cadence, sleep power, sensing, communication, and the rest of the device.

The energy gap between paradigms is not a matter of degree but of scale, spanning the eight orders of magnitude sketched in figure 2.6. Table 2.12 grounds that span in concrete numbers, translating each paradigm’s energy cost into how many inferences a single smartphone battery could sustain.

³² | **Microcontroller (MCU):**

A single-chip computer whose design prioritizes minimal cost and power over performance, creating the “radical constraint” mentioned. Many TinyML targets provide only tens to hundreds of kilobytes of on-chip SRAM and no virtual memory; some higher-end microcontrollers can attach external memory. The complete model, activations, and runtime state must therefore fit the memory actually available to the target, often orders of magnitude less than a smartphone provides.

³³ | **TinyML energy gap:**

This differential is rooted in hardware design philosophy; cloud GPUs are optimized for raw throughput, consuming hundreds of watts, while TinyML microcontrollers are designed for near-zero power sleep states. For ordinary cloud inference, a single request may consume about 10 joules while a specialized TinyML device uses about 10 microjoules, a 1,000,000× gap. Cloud LLM queries can push the comparison even further, as quantified in table 2.12.

³⁴ | **Coin-cell deployment:**

A CR2032 battery (225 mAh at 3 V, ~675 mWh) ideally powers a complete 10–50 μ W system for about 1.5–7.7 years. This “deploy-and-forget” operating model drives innovation in intermittent computing, where the device sleeps between inferences to stretch the energy budget across years of unattended operation.

Table 2.12: Energy per inference across paradigms: Analytical device-energy estimates and scale anchors, with the corresponding inference counts from one smartphone battery. These are not controlled full-system measurements.

Paradigm	Example Workload	Energy/Inference	Inferences per Battery (3.7 V, 3000 mAh)
Cloud	GPT-4 query	~1 kJ	40 queries
Cloud	ResNet-50 (accelerator server)	75.5 mJ	529,505 queries
Edge	ResNet-50 (Jetson)	15.9 mJ	2,511,628 queries
Mobile	MobileNet (NPU)	2.60 mJ	15,393,505 queries
TinyML	Keyword spotting	10 μ J	3996 million queries

Operating on sub-milliwatt power budgets enables continuous sensing without grid power, but forces the complete inference state into on-chip SRAM. The four-column taxonomy in figure 2.9 traces how kilobyte memory and duty-cycled sleep states enable always-on workloads while restricting model updates.




Figure 2.9: TinyML Constraint Map: Milliwatt power and kilobyte memory make always-on sensing economically possible, but those same limits force the model, activations, runtime, and update path to fit a radically smaller engineering envelope.

2.8.1 TinyML advantages and operational trade-offs

TinyML operates at hardware extremes. Compared to cloud systems, TinyML deployments commonly provide roughly 10^6 to 10^9 times less memory, depending on whether the microcontroller budget is in low megabytes or kilobytes, with power budgets in the milliwatt range. These strict limitations enable months or years of autonomous operation³⁵ but demand specialized algorithms, model compression, and careful systems co-design. Devices range from palm-sized developer kits to millimeter-scale chips,³⁶ enabling ubiquitous sensing in contexts where networking, power, or maintenance are costly. Representative developer kits include the Arduino Nano 33 BLE Sense (256 KB RAM, 1 MB flash, 20–40 mW) and ESP32-CAM (520 KB RAM, 4 MB flash, 50–250 mW).

TinyML's extreme resource constraints paradoxically enable unique advantages. By avoiding network transmission entirely, TinyML devices avoid network-transmission latency, enabling rapid local responses for sensing and control loops without communication overhead. This self-sufficiency also transforms the economics of large-scale deployments: when per-node costs drop to single-digit dollars, instrumenting an entire factory floor, farm, or building becomes financially viable in ways that edge or cloud alternatives cannot match. Energy efficiency compounds the economic case, enabling multi-year operation on small batteries or even indefinite operation through energy harvesting. Privacy benefits follow naturally from locality, because raw data never leaves the device, reducing transmission risks and simplifying compliance. On-device processing alone does not automatically provide formal privacy guarantees without additional security mechanisms.

³⁵ | **On-device training constraints:** Full on-device training must keep intermediate layer outputs so later weight updates can reuse them, consuming memory proportional to model depth; during inference, each layer's temporary values can usually be discarded after the next layer consumes them; during training, many of those values must remain available so the update step can decide how weights should change. With only 256 KB–2 MB RAM, full training generally exceeds this memory budget; TinyTL freezes weights and updates biases to reduce activation memory (Cai et al. 2020). This memory constraint is why TinyML devices are predominantly inference-only, with model updates pushed via firmware rather than learned in situ.

³⁶ | **TinyML device range:** This physical range reflects a direct trade-off between deployment context and computational capability. Millimeter-scale systems prioritize minimal power (~140 μ W) for single-function, long-duration tasks, whereas palm-sized boards trade larger size and higher power for the ability to process multiple complex sensor streams. This co-design choice creates a $> 10,000\times$ power and $\sim 100\times$ area difference across the operational spectrum of TinyML devices.

These capabilities require substantial trade-offs. Computational constraints impose severe limits: microcontrollers commonly provide 10^5 to 10^6 bytes of RAM, forcing models and intermediate activations into the tens-of-kilobytes to low-megabytes range depending on the workload. Development complexity requires expertise spanning neural network optimization, hardware-level memory management, embedded toolchains, and specialized debugging across diverse microcontroller architectures.

Beyond these technical constraints, operational challenges compound the difficulty. An energy-harvesting device can satisfy its long-term average power budget yet fail when stored energy is exhausted before the next harvest. Its duty cycle must therefore respect both energy per inference and the timing of available power. Model quality can suffer from aggressive compression and reduced precision, limiting suitability for applications requiring high accuracy or robustness. Deployment can also be inflexible: devices may run a small set of fixed models, and updates may require firmware workflows that are slower and riskier than cloud rollouts. Ecosystem fragmentation³⁷ across microcontroller vendors and ML frameworks creates additional overhead and portability challenges.

2.8.2 Environmental and health monitoring

TinyML applications belong here when ultra-low power, low per-node cost, and local processing make a deployment feasible that no other paradigm can sustain. Wake-word detection is the most familiar consumer example: a purpose-built low-power component listens continuously, processes audio locally, and activates higher-power components only when a wake phrase is detected, reducing average power draw.³⁸ Precision agriculture exploits the same locality pressure from a different direction: FarmBeats used sensors, cameras, drones, and local gateway processing to reduce raw data movement where farm connectivity was costly (Vasisht et al. 2017). TinyML pushes that locality logic further when the sensing node itself must run under milliwatt budgets.

Wildlife conservation uses TinyML for remote environmental monitoring, where solar-powered audio sensors can process continuous streams locally for species identification. Under this chapter's illustrative payload assumptions, local analysis reduces transmission from 4.3 GB/day of raw audio to 400 KB/day of detection summaries, a 10,750× reduction. Medical wearables apply the same local-processing logic to health. Sensitivity, sampling rate, power, and battery life are device-specific validation targets rather than universal properties of TinyML hardware. Local processing can reduce continuous transmission and support privacy, while total runtime still depends on sensors, radios, duty cycle, and sleep power.

The four deployment paradigms now span the full range from megawatt data centers to milliwatt microcontrollers. Each paradigm emerged as a response to specific physical constraints, and each excels within its operating envelope. The question of how an engineer should choose among them, and what to do when no single paradigm satisfies all requirements, motivates the comparative analysis that follows.

2.9 Paradigm Selection

An architect choosing where to run an ML feature rarely optimizes one dimension. A privacy rule may forbid cloud processing, a latency budget may forbid remote inference, a memory footprint may exceed the mobile device, and a cost target may rule out always-on edge hardware. Cloud, edge, mobile, and TinyML are operating envelopes for those conflicts, so selecting among them requires a unified comparison framework and a structured decision process.

2.9.1 Comparative trade-off analysis

Deployment decisions require seeing latency-vs-throughput paradigm trade-offs side by side across the dimensions that matter. A system architect choosing between edge and mobile deployment must compare latency, power, cost, privacy, and development complexity simultaneously. Table 2.13 provides this comparison across fourteen dimensions, from compute power and latency to cost and deployment speed.

³⁷ | **TinyML ecosystem fragmentation:** TinyML spans microcontroller families such as Arm Cortex-M, RISC-V, and Xtensa, each with different instruction sets, memory layouts, kernels, and vendor toolchains; cloud and mobile stacks also vary, but their runtimes usually hide more of this hardware diversity. A model optimized for one microcontroller target may therefore require retuning and revalidation for another, multiplying the engineering cost of multi-device deployment.

³⁸ | **Always-on wake-word detection:** A small, specialized model listens for the acoustic signature of a wake phrase and acts as a power gate for the higher-power application processor. The acceptable detector budget depends on the platform, but spending even a few extra milliwatts continuously can materially shorten battery life. The architecture succeeds only when the always-on path consumes far less average power than the components it prevents from waking.

Table 2.13: Fourteen-Dimension Paradigm Comparison: A comprehensive side-by-side comparison across fourteen dimensions that matter for deployment decisions. Cloud ML provides the strongest compute, while local processing can reduce transmission and centralized-storage exposure without itself guaranteeing privacy. This table serves as the primary reference for system architects evaluating deployment options.

Aspect	Cloud ML	Edge ML	Mobile ML	TinyML
Processing Location	Centralized cloud servers (Data Centers)	Local edge devices (gateways, servers)	Smartphones and tablets	Ultra-low-power microcontrollers and embedded systems
Latency	100–500 ms	10–100 ms	5–50 ms	1–10 ms
Compute Power	Very High (Multiple GPUs/TPUs)	High (Edge GPUs)	Moderate (Mobile NPUs/GPUs)	Very Low (MCU/tiny processors)
Storage Capacity	Cloud-scale (petabytes+)	Large (terabytes)	Moderate (gigabytes)	Very Limited (kilobytes–megabytes)
Energy Consumption	Very High (kW–MW range)	High (hundreds of watts)	Moderate (3–5 W)	Very Low (mW range)
Scalability	High (provider capacity limits)	Good (limited by edge hardware)	Moderate (per-device scaling)	Limited (fixed hardware)
Data Exposure	Remote processing expands the surface	Local processing reduces transmission	On-device processing reduces transmission	Raw data can remain local
Connectivity Required	Constant high-bandwidth	Intermittent	Optional	None
Offline Capability	None	Good	Excellent	Complete
Real-time Processing	Dependent on network	Good	Very Good	Excellent
Cost	High (\$1000s+/month)	Moderate (\$100s–\$1000s)	Device purchase (\$200–\$1000+; often user-owned)	Very Low (\$1–\$10)
Hardware Requirements	Cloud infrastructure	Edge servers/gateways	Modern smartphones	MCUs/embedded systems
Development Complexity	High (cloud expertise needed)	Moderate-High (edge+networking)	Moderate (mobile SDKs)	High (embedded expertise)
Deployment Speed	Fast	Moderate	Fast	Slow

This contrast reflects a common trade-off between data locality and computational scale. Remote processing increases the disclosure surface and requires additional safeguards; local processing reduces transmission and centralized-storage exposure but does not itself guarantee privacy. The archetype-paradigm mapping established in section 2.3 connects these characteristics to specific workload requirements, with each archetype gravitating toward paradigms that address its binding constraint.

No single paradigm dominates across all engineering axes. In figure 2.10, four-axis polar radar charts evaluate the paradigms on an ordinal 0–10 scale: subplot (a) maps raw compute and scalability against latency and energy efficiency, while subplot (b) evaluates operational autonomy across connectivity, locality, real-time response, and offline execution.

The radar plots make the deployment decision explicit: no paradigm dominates all axes. Each paradigm concentrates complexity in a different place: cloud requires distributed-infrastructure and cost-governance expertise; edge and TinyML add fleet management, hardware diversity, and embedded constraints; mobile adds platform-specific integration and release processes. Cost structures likewise differ rather than following one ranking: cloud emphasizes recurring service expense, edge emphasizes site hardware and operations, mobile often uses user-owned devices, and TinyML trades low unit cost for specialized development and large-fleet maintenance.

A critical pitfall in deployment selection is choosing paradigms based solely on model accuracy without considering system-level constraints. A cloud-deployed model achieving 99 percent accuracy becomes useless for autonomous emergency braking if network latency exceeds reaction time requirements; a high-accuracy edge model that drains a mobile device’s battery in minutes fails




Figure 2.10: Paradigm Comparison Radar Plots: Two radar plots compare performance and operational characteristics across cloud, edge, mobile, and TinyML paradigms using ordinal 0–10 scores. The left plot contrasts compute power, low-latency capability, scalability, and energy efficiency; the right plot contrasts connectivity independence, data locality, real-time capability, and offline operation. In both plots, a larger polygon indicates stronger performance, with cloud ML peaking on compute and scalability and TinyML peaking on energy efficiency, data locality, and offline operation.

despite superior accuracy. Successful deployment requires evaluating latency requirements, power budgets, network reliability, data privacy regulations, and total cost of ownership simultaneously. These constraints should be established before model development to avoid expensive architectural pivots late in the project.

2.9.2 Decision framework

Selecting the appropriate deployment paradigm requires a decision framework based on application constraints rather than organizational biases or technology trends. The gates are ordered by how quickly they can invalidate an architecture: privacy can make remote processing impermissible, latency can make it physically impossible, compute demand can make a local target infeasible, and cost compares the feasible survivors as operational architectures. Use figure 2.11 as a screening aid that organizes privacy, latency, compute-demand, and cost questions before a final architecture is selected.

The framework groups four critical decision layers rather than computing a unique answer. Privacy determines whether remote processing is permissible. Latency rejects paths that miss response-time budgets, compute demand rejects targets that cannot run the workload, and cost compares the feasible survivors. Because the drawn branches reconverge and the low-cost branch retains both mobile ML and TinyML, figure 2.11 narrows the design space but does not select a single paradigm.

A safety-critical braking example makes the sequencing concrete, because latency can eliminate cloud deployment before compute or cost are considered.

Example 2.1: Autonomous vehicle emergency braking

Scenario: An autonomous vehicle vision system performs real-time pedestrian detection for emergency braking under a strict 100 ms end-to-end latency budget.

Diagnosis: The scenario assumes 50–150 ms of network round-trip delay for cloud offload and 10–30 ms for local object detection on an automotive edge accelerator (300 GFLOP/s). Under those assumptions, the remote path alone causes the vehicle to travel 2.8 m before braking.

Systems lesson: Hard real-time safety constraints eliminate cloud offloading regardless of compute availability. Latency bounds dictate an Edge ML deployment architecture, reserving cloud connections for offline model training and telemetry aggregation.




Figure 2.11: Deployment Decision Screening: Four layers organize privacy, latency, compute-demand, and cost questions, then associate their branches with cloud, edge, mobile, and TinyML options. Reconverging branches and a low-cost branch that reaches both mobile and TinyML make this a screening aid rather than a deterministic selector.

This decision framework identifies technically feasible options, but feasibility does not guarantee success. Production deployment also depends on organizational capabilities that determine whether a technically sound choice can be implemented and maintained effectively.

Successful deployment requires considering factors beyond pure engineering constraints. Team expertise must align with paradigm requirements: cloud ML demands distributed systems knowledge, edge ML requires device management capabilities, mobile ML needs platform-specific optimization skills, and TinyML requires embedded systems expertise. Organizations lacking appropriate skills face extended development timelines that can undermine even the strongest technical advantages. Monitoring and maintenance capabilities similarly determine viability at scale: edge deployments require distributed device orchestration, while TinyML demands specialized firmware management that many organizations lack. Cost structures add another dimension, because the temporal pattern of expenses varies dramatically across paradigms. Cloud incurs recurring operational costs favorable for unpredictable workloads; Edge requires substantial upfront investment offset by lower ongoing costs; Mobile uses user-provided devices to minimize infrastructure expenses; and TinyML minimizes hardware costs while demanding significant development investment.

These organizational realities surface a broader concern: a machine learning approach is not always the right choice. Every ML deployment carries operational overhead (including data pipelines, monitoring, and retraining infrastructure) that simpler heuristic systems avoid, and that overhead must be justified by measurably better outcomes.

Systems Perspective 2.4: The complexity tax

Before we commit to any ML deployment, we must weigh the operational burden against simpler alternatives.

Consider a classification problem solvable by either a heuristic (if-then rules) or a deep learning pipeline. The heuristic may be fifty lines of code with near-zero compute cost, about one hour per month to update rules, and no model drift. The ML system may still have only fifty lines of model code, but it also brings roughly 2,000 lines of infrastructure for data pipelines, monitoring, and GPU drivers, plus about 40 hours per month debugging drift and managing infrastructure.

An ML system that improves accuracy from 90 percent to 95 percent may still be a poor engineering choice if it introduces a $40\times$ increase in complexity. ML systems engineering is the art of minimizing this tax through robust architecture. If the operational cost of maintaining model quality over time is unaffordable, the simpler heuristic may be the superior systems choice.

That comparison turns model choice into a system-design decision spanning physics, infrastructure cost, and the ongoing burden of maintaining accuracy. When the **complexity tax** exceeds the accuracy gain, the simpler heuristic is the superior systems choice.

Checkpoint 2.2: System design

The central trade-off is often accuracy vs. complexity.

Decision gates

- ☐ **The baseline:** have you measured the accuracy of a simple heuristic (regex, logistic regression) before training a deep network?
- ☐ **The infrastructure cost:** is the 2 percent accuracy gain from a transformer worth the $10\times$ inference cost and maintenance burden compared to a smaller model?

Successful deployment balances technical optimization against organizational capability. Paradigm selection extends well beyond technical requirements to encompass team skills, operational capacity, and economic constraints, all bounded by the physical scaling laws developed in this chapter. Chapter 14 and chapter 12 develop the operational and measurement consequences. In practice, however, the decision framework rarely points to a single winner. Most production systems combine multiple paradigms, such as training in the cloud, serving at the edge, and preprocessing on mobile, to satisfy constraints that no single deployment target can meet alone.

2.10 Hybrid Architectures

The decision framework (figure 2.11) narrows the feasible paradigms for a given application. In practice, production systems rarely use just one paradigm. Voice assistants combine TinyML wake-word detection with mobile speech recognition and cloud natural language understanding. Autonomous vehicles pair edge inference for real-time perception with cloud training for model updates. These hybrid architectures exploit the strengths of multiple paradigms while mitigating their individual weaknesses. Three integration strategies formalize how such combinations work in practice.

2.10.1 Integration patterns

The three essential hybrid ML patterns differ by which boundary creates the constraint: train-serve split, hierarchical processing, or progressive deployment. Their selection is an iron-law decision: each stage should run where its binding term, whether training compute, local latency, or model size, is cheapest to satisfy.

The **train-serve split** places training in the cloud while inference happens on edge, mobile, or tiny devices. This pattern exploits cloud scale for training while benefiting from local inference

latency and data locality. When a trained artifact serves many requests, its fixed training cost can be amortized across repeated inference.³⁹

³⁹ | **Train-serve cost asymmetry:** Training usually requires many forward and backward passes, whereas one inference request uses a trained artifact for one or more forward passes. Training is not literally one-time because models may be re-trained, and inference cost varies sharply with model size, sequence length, batching, and service overhead. The economic rationale for the split is amortization: repeated use can spread each training cycle's fixed cost across many served predictions.



Definition 2.6: Hybrid ML

Hybrid machine learning is the deployment strategy that distributes an ML pipeline across two or more deployment tiers, assigning each stage according to its latency, compute, data, and operational constraints.

1. **Significance:** Hybrid architectures exploit the iron law's additive structure: the edge tier can reduce L_{lat} for time-sensitive preprocessing and inference, while the cloud tier can provide the R_{peak} needed for training, retraining, and heavy batch inference. The split follows the data locality invariant: compare each stage's complete local and remote paths, reject paths that miss the service deadline, and then choose among the feasible alternatives.
2. **Distinction:** Unlike cloud-only deployment (which accepts the distance penalty for all stages) or edge-only deployment (which accepts limited R_{peak} for all stages), hybrid ML dynamically assigns each pipeline stage to the tier where its binding iron law term is minimized.
3. **Common pitfall:** A frequent misconception is that hybrid ML is just "running two models." In reality, the participating tiers must share synchronized state—feature definitions, model versions, and preprocessing logic—so that their paths produce consistent results. Without this synchronization, training-serving skew emerges at tier boundaries.

In **hierarchical processing**, data and intelligence flow between computational tiers. TinyML sensors perform basic anomaly detection, edge devices aggregate and analyze data from multiple sensors, and cloud systems handle complex analytics and model updates. Each tier handles tasks appropriate to its capabilities.

The third pattern, **progressive deployment**, systematically compresses models for deployment across tiers. A large cloud model becomes progressively optimized versions for edge servers, mobile devices, and tiny sensors. Voice assistants exemplify this pattern: wake-word detection uses small, always-on models, often tens of kilobytes for benchmark TinyML neural networks and sub-milliwatt to milliwatt-scale power on dedicated low-power hardware, while complex natural language understanding requires much larger models in cloud infrastructure.

With three integration patterns available, selection becomes a constraint-matching problem: choose the pattern whose trade-off profile matches the system's dominant bottleneck. Table 2.14 summarizes the trade-off, the conditions that favor each pattern, and the conditions that argue against it.

Table 2.14: Hybrid Pattern Selection Guide: The trade-off, favoring conditions, and disqualifying conditions for each of the three hybrid integration patterns. Selection matches the pattern's trade-off profile to the system's dominant bottleneck.

Pattern	Trade-off	Choose when	Avoid when
Train-Serve Split	Training cost vs. inference latency	Training requires scale that inference does not; privacy matters for inference but not training	Model needs continuous learning from deployed data
Hierarchical Processing	Local autonomy vs. global optimization	Data volume exceeds transmission capacity; decisions needed at multiple timescales	All processing can occur at one tier; network is reliable and fast
Progressive Deployment	Model quality vs. deployment reach	Same model needed at multiple capability levels; graceful degradation required	Model cannot be meaningfully compressed; single deployment target

Voice assistants combine train-serve split, progressive deployment, and hierarchical processing; autonomous vehicles combine hierarchical processing with progressive deployment to run optimized

models at each tier. Privacy-preserving distributed training approaches extend this menu when data should remain close to the devices that produce it.

2.10.2 Production system integration

Production hybrid systems coordinate multiple tiers through structured interaction channels. In figure 2.12, labeled connection vectors illustrate this multi-tier architecture, showing how trained model artifacts deploy downward while sensor streams and intermediate inference results propagate upward to centralized analytics.




Figure 2.12: Hybrid System Interactions: Data flows upward from sensors through processing layers to cloud analytics, while trained models deploy downward to edge, mobile, and TinyML inference points. Five connection types (deploy, data, results, assist, and sync) establish a distributed architecture where each paradigm contributes unique capabilities.

Production systems demonstrate these integration patterns by placing each tier boundary at a different binding constraint. Industrial defect detection exemplifies Train-Serve Split: cloud infrastructure trains vision models on datasets from multiple facilities, then distributes optimized versions to edge servers managing factory floors, tablets for quality inspectors, and embedded cameras on production lines. Agricultural monitoring illustrates Hierarchical Processing: soil sensors perform local anomaly detection at the TinyML tier, edge processors aggregate data from dozens of sensors and identify field-level patterns, while cloud infrastructure handles farm-wide analytics and seasonal planning. Fitness tracking exemplifies Progressive Deployment with gateway patterns: wearables continuously monitor activity using microcontroller-optimized algorithms consuming < 1 mW, sync processed summaries to smartphones that combine metrics from multiple sources, then transmit periodic updates to cloud infrastructure for longitudinal health analysis.

2.10.3 Why hybrid approaches work

The four deployment paradigms differ in physical resource budgets but share systems engineering foundations. In figure 2.13, a three-tier vertical hierarchy traces how top-level deployment targets depend on shared data, compute, and architectural principles, which in turn support cross-cutting efficiency and reliability considerations.

This convergence explains why techniques transfer between scales when they attack a shared bottleneck. Cloud-trained models can deploy to edge because the learned weights and operator graph can be reused, but the target device changes the memory, precision, latency, and power budget. Lower-precision representations developed for edge deployment reduce cloud serving costs; chapter 10 formalizes these methods as quantization. Strategies for splitting work across devices likewise inform edge deployments that partition one model across more than one processor; chapter 8 formalizes this family as model parallelism.

Mobile optimization insights inform cloud efficiency because memory bandwidth constraints appear at every scale. Methods that reduce memory traffic on phones can also reduce cloud inference costs when applied to batch serving. TinyML’s extreme constraints likewise motivate

compact representations and memory-management techniques that can inform larger systems under similar memory pressure.




Figure 2.13: Convergence of ML Systems: Despite vast resource disparities across cloud, edge, mobile, and TinyML tiers, all deployments converge on three shared systems foundations: data pipelines, resource management, and hardware-software architecture. These shared pillars in turn support unified requirements for performance optimization, operational monitoring, and trustworthy engineering.

The same layered pattern continues through chapter 4 for data pipelines, chapter 10 for optimization, and chapter 14 for operational aspects. All of these apply whether the target is a TPU Pod or an ESP32. Shared principles also mean shared vulnerabilities: data drift, model decay, and monitoring recur at every tier, so the remaining lessons must account for them.

Checkpoint 2.3: Hybrid ML patterns

Hybrid architectures work when partitioning work across tiers, not when copying the same pipeline everywhere.

Integration patterns

- ☐ **Train-serve split:** can you explain why training in the cloud and serving on edge/mobile is often economically optimal, even when the model runs locally?
- ☐ **Hierarchical processing:** can you describe what each tier does in a sensor → edge → cloud pipeline, and why pushing some decisions down reduces both latency and bandwidth?
- ☐ **Progressive deployment:** can you explain how one model family becomes multiple deployed artifacts (cloud, edge, mobile, tiny) through systematic compression?

Design sanity checks

- **Boundary choice:** given a concrete application, can you justify where the tier boundary should fall (latency, privacy, bandwidth, power), not just what model to use?
- **State synchronization:** can you identify which feature definitions, model versions, and preprocessing logic must remain synchronized across tiers?

2.11 System Entropy: Why Deployment Is Not the End

The shared foundations in figure 2.13 also share a vulnerability. Deployment is not the end of the engineering challenge; it is the beginning of a new one. A sorting algorithm retains its specified behavior while its code, execution environment, and input contract remain valid. ML systems face an additional form of **system entropy**: statistical decay caused by the gap between training conditions and live operating conditions.

An ML model's accuracy can change when the world drifts away from its training distribution, even if its code and weights remain unchanged. Equation 2.8 provides a local diagnostic approximation: system quality may decline as the distance between the training distribution and the live data distribution grows, at a rate proportional to the model's sensitivity to distributional shift:

$$\Delta\text{Quality} \approx -S_{\text{shift}} \cdot \mathcal{D}(P_{\text{train}}, P_{\text{live}}) \quad (2.8)$$

Reliability therefore requires monitoring to determine whether updates are warranted. The operational aspects covered in chapter 14 address precisely this challenge.

 War Story 2.1: The Zillow Offers collapse (2021)

Context: Zillow, a real-estate marketplace, launched “Zillow Offers” to buy homes directly through an iBuying business that depended on forecasting home prices and resale economics (Zillow Group 2021).

Mechanism: Zillow reported that it had purchased homes at prices above its revised estimates of future selling prices, while resale-capacity constraints delayed some closings (Zillow Group 2021). Forecast uncertainty became inventory exposure because the system coupled price estimates to purchase commitments.

Impact: Zillow wrote down \$304 million in housing inventory, laid off 25 percent of its workforce (2,000 employees), and permanently shut down the Offers division.

Response: A safer design would cap inventory exposure, preserve manual overrides, and place circuit breakers between forecast uncertainty and purchase decisions.

Systems lesson: Distribution shift is not just a metric drop; it is a business risk. Automated decision-making systems interacting with dynamic markets require rapid feedback loops and circuit breakers, not just accurate offline models.

Zillow's collapse is not merely a cautionary tale. The failure involved forecast uncertainty and the system that acted on it: automated buying coupled uncertain estimates to inventory and balance-sheet commitments. ML systems engineering makes that propagation path visible before losses compound.

2.12 Fallacies and Pitfalls

Beyond statistical decay, engineers also fall prey to common misconceptions about ML deployment. The physical constraints examined throughout this chapter create counterintuitive behaviors that challenge intuitions from traditional software engineering. These fallacies and pitfalls capture architectural mistakes that waste development resources, miss performance targets, or deploy systems critically mismatched to their operating constraints.

Fallacy: *One deployment paradigm solves all ML problems.*

Physical constraints create hard boundaries that no single paradigm can span. The memory-wall discussion in section 2.4 shows that bandwidth, memory capacity, and latency scale differently from raw compute, producing qualitatively different bottlenecks across paradigms. Table 2.13 lists illustrative deployment envelopes rather than physics-only latency bounds. A real-time robotics system requiring sub-10 ms response cannot use a remote path whose end-to-end latency exceeds that budget, and a billion-parameter language model *cannot* fit on a microcontroller with 256 KB RAM regardless of model-size reduction. The optimal architecture typically combines paradigms, such as cloud training with edge inference or mobile preprocessing with cloud analysis.

A related misconception holds that moving computation closer to the user always reduces latency, ignoring the processing overhead introduced by less powerful edge hardware, a trade-off explored in inference benchmarks (section 12.8).

Pitfall: *Relying on model optimization to overcome mobile power and thermal limits.*

Compression techniques do not scale indefinitely against physics. Consider a smartphone with a 15 Wh battery. A light inference workload drawing 1 W runs for $\frac{15Wh}{1W} = 15$ h, but a heavy workload drawing 5 W, common for large on-device models, drains the same battery in $\frac{15Wh}{5W} = 3$ h.

The 5 W workload can also trigger thermal throttling. This scenario budgets approximately 3 W for sustained passive cooling. A 4× reduction is an illustrative redesign target, not a universal consequence of quantization. A continuous workload that still exceeds the measured thermal envelope needs a lower duty cycle, a different model, or different hardware.

Fallacy: *TinyML represents scaled-down mobile ML.*


The difference is qualitative, not just quantitative. As section 2.8.1 establishes, TinyML microcontrollers provide 256 KB to 1 MB of memory vs. mobile devices with 8–16 GB, a 10,000× difference requiring entirely different algorithms. Both mobile ML and TinyML use reduced-precision arithmetic; accuracy effects depend on the model, dataset, method, and hardware. Mobile devices run models with millions of parameters; TinyML models often contain tens to hundreds of thousands of parameters, demanding distinct architectural choices such as specialized lightweight operations designed to minimize multiply-accumulate counts. Power budgets show similar discontinuities: mobile inference consumes 3–5 W, while TinyML targets 1–10 mW for battery-free energy harvesting. These thousand-fold gaps make TinyML a distinct problem class, not a smaller version of mobile ML. Teams that apply mobile optimization techniques directly to TinyML projects discover that quantization from FP32 to INT8 is insufficient when models must fit in 64 KB, forcing complete architectural redesign.

Pitfall: *Minimizing computational resources minimizes total cost.*

Teams optimize per-unit resource consumption while ignoring operational overhead and development velocity. As the decision framework in section 2.9.2 emphasizes, paradigm selection requires evaluating total cost of ownership, not just compute costs. In this illustrative scenario, a cloud inference service costing \$2,000/month in compute appears expensive vs. \$500/month in edge-hardware amortization, but edge deployment adds network engineering (\$3,000/month), hardware maintenance (\$500/month), and reliability engineering (\$2,000/month), totaling \$6,000/month—a 3× difference. If the cloud version reaches production in two months and custom edge infrastructure takes six, the four-month delay is another scenario cost rather than a universal schedule. The optimal cost solution requires total cost of ownership analysis including development time, operational complexity, and opportunity costs, not merely minimizing compute expenses.

Fallacy: *Model optimization translates linearly to system speedup.*

Amdahl's Law⁴⁰ establishes hard limits that the bottleneck principle (section 2.3.1) makes operational, where section D.2.3.1 derives the strong-scaling form and works a speedup example at eight processors: $\text{Speedup}_{\text{overall}} = \frac{1}{(1-p) + \frac{p}{s}}$ where p is the fraction of work that can be improved and s is the speedup of that fraction. Consider tapping the shutter on a smartphone camera. The image passes through 100 ms of signal processing (auto-exposure, white balance), 60 ms of ML scene classification, and 40 ms of postprocessing (tone mapping, HDR merge)—200 ms total. Optimizing the ML classifier to run 10× faster (6 ms instead of 60 ms) drops total time from 200 ms to 146 ms—only



A faster model stage does not linearly speed up a camera pipeline.

⁴⁰ **Amdahl's law:** Formalized by Amdahl (1967) for multiprocessor scaling, this principle applies directly to ML deployment pipelines where the model is only one stage among many. In this camera pipeline example, ML inference is 60 ms of 200 ms; even a 100× model speedup yields only about 1.4–2× end-to-end improvement because the rest of the pipeline is unchanged. Teams that benchmark model latency in isolation systematically overestimate deployment gains.

1.37× overall, not 10×. Even eliminating ML entirely ($s = \infty$) achieves only 1.43× speedup, because the remaining 70 percent of the pipeline is untouched. Effective optimization requires profiling the entire pipeline and addressing bottlenecks systematically, because system performance depends on the slowest unoptimized stage.

Pitfall: *Assuming more training data always improves deployed model performance.*

Three constraints limit data scaling benefits, as the workload archetypes in section 2.3 illustrate. First, model capacity can limit what additional data contributes. Second, data quality can outweigh quantity because mislabeled examples and misleading patterns degrade performance. Third, deployment distribution matters: in-domain data can outperform much larger out-of-domain datasets.

Fallacy: *One model binary can serve every edge hardware target efficiently.*

Teams build a single model artifact and deploy it identically to every target device, treating deployment as a packaging step rather than an optimization opportunity. In practice, hardware-specific optimizations can yield substantial efficiency gains that generic binaries cannot capture. An INT8 model running on a device with a dedicated Neural Processing Unit (NPU) can achieve higher throughput per watt than the same model running in FP32 on a general-purpose CPU, because the NPU's fixed-function INT8 datapaths avoid the energy overhead of floating-point arithmetic. Similarly, operator fusion (combining adjacent operations so intermediate tensors are not written back to memory) and memory layout tuning for a specific accelerator's cache hierarchy can reduce inference latency without changing the model's weights. As the deployment paradigm analysis in section 2.1 establishes, each paradigm imposes distinct hardware constraints; a model binary optimized for an Arm Cortex-A78 will underutilize the matrix acceleration units on a device equipped with an Arm Ethos-U NPU. Teams that skip per-target optimization either waste battery life on mobile devices or fail to meet latency SLAs on edge hardware, forcing costly postdeployment remediation. A generic artifact can remain a portability baseline, but profiling on the production runtime, delegate, memory hierarchy, and thermal envelope must determine whether it is acceptable, inefficient, or fundamentally mismatched.

2.13 Summary

ML systems engineering is governed by the physical reality that every deployment target enforces a distinct hardware contract. The iron law of ML systems ($T = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}}$) unifies these physical limits into a single execution model: propagation delay contributes fixed latency (L_{lat}), the memory wall bounds data movement (D_{vol}/BW), and thermal power budgets cap sustained compute. Across cloud, edge, mobile, and TinyML tiers, the categorical power envelope from megawatts to milliwatts spans nine orders of magnitude, while representative systems occupy narrower ranges within it.

By applying the bottleneck principle to workload archetypes, engineers determine whether a model's latency is dominated by compute efficiency, memory bandwidth, or network transit before committing to single-node or hybrid architectures. The quantitative trade-offs (table 2.13) and systematic decision framework (figure 2.11) developed in this chapter establish the physical baselines that govern all downstream decisions in model compression, data pipeline design, and high-performance serving.



Key Takeaways: Same model, different engineering

- **Physical constraints define feasibility:** Speed of light (~36 ms cross-country round-trip), power wall, memory wall, and latency budgets create hard boundaries that engineering cannot overcome, only navigate.
- **Identify bottlenecks before optimizing:** The same model is compute bound in training but memory bound in inference. The iron law and bottleneck principle pinpoint which constraint dominates; optimizing the wrong term yields zero speedup.
- **Workload archetypes determine deployment feasibility:** A Compute Beast such as large-scale image-model training often needs cloud-scale resources; a Tiny Constraint such as

keyword spotting requires microcontroller efficiency. The same optimization strategy cannot serve both; engineers must match the workload archetype to the deployment paradigm.

- **Deployment power spans nine orders:** Cloud data-center infrastructure operates at megawatt scale, while TinyML systems target milliwatts. This gap enables entirely different application classes rather than representing a limitation.
- **Hybrid architectures are prevalent in production systems:** Voice assistants span TinyML (wake-word), mobile (speech-to-text), and cloud (language understanding). When one paradigm does not suffice, integration patterns (train-serve split, hierarchical processing, progressive deployment) formalize how paradigms combine.
- **System-level speedup obeys Amdahl's Law, not model-level gains:** In the worked camera pipeline, a $10\times$ faster model yields only $1.37\times$ system speedup when ML accounts for 30 percent of the pipeline. Profile the full system before optimizing any component.
- **Deployment begins a feedback loop:** Live data can diverge from training conditions while code and weights stay fixed, so monitoring and bounded operational exposure are part of the system architecture.


Deployment paradigms are not arbitrary design preferences; they are continuous operational regions dictated by physical laws. Selecting a deployment target changes which term of the iron law is most likely to bind: data-center training can push against peak compute (O/R_{peak}), mobile inference can collide with memory bandwidth (D_{vol}/BW), and remote real-time inference can be constrained by propagation delay (L_{lat}). Systems engineering begins by measuring that binding constraint. Optimizing model architecture or software runtimes without respecting the underlying hardware contract produces fragile systems that fail under real-world operational pressure.



What's Next: From theory to process

While the physical constraints and decision framework established in section 2.9.2 define where an ML system can physically execute, physical feasibility alone does not guarantee operational success over time. As demonstrated by Zillow's \$304 million write-down (Zillow Group 2021), deploying models into production requires containing forecast uncertainty and managing operational risk across the software lifecycle. Chapter 3 introduces the ML development lifecycle, the systematic engineering process that carries a system from problem formulation through deployment, monitoring, and maintenance.

ML Workflow




- 3.1 ML Lifecycle
- 3.2 Lifecycle Stages
- 3.3 Problem Definition
- 3.4 Data Collection
- 3.5 Model Development
- 3.6 Evaluation and Validation
- 3.7 Deployment and Integration
- 3.8 Monitoring and Maintenance
- 3.9 Systems Thinking
- 3.10 Fallacies and Pitfalls
- 3.11 Summary

Purpose

Why is seeing the whole map necessary before walking any single path?

The D·A·M taxonomy names the components of every ML system, and deployment location determines the physical constraints each component must satisfy. Teams often treat these as separate concerns, with one team collecting data, another designing the model, and a third provisioning hardware. Yet the taxonomy's deepest lesson is that these components interact. The collected data constrains which algorithms are feasible. The chosen algorithm dictates what hardware can run it. The target hardware reshapes what data can be processed. Pull on any single thread and the entire system shifts. These interactions play out across components and across time. A model that performs well at launch may degrade as the data distribution shifts, prompting investigation and, when evidence warrants it, model or data updates. Optimizing each piece in isolation is how teams build accurate models that cannot be deployed and efficient pipelines that feed the wrong data. A data engineer who sees how preprocessing choices constrain downstream architectures builds different pipelines than one who treats data preparation as an isolated task; a model developer who knows the deployment target's memory budget from day one makes different architecture decisions than one chasing accuracy in a vacuum. Before the details of any one component can be understood, the full map must show how an ML system is built, evaluated, and sustained as a coherent whole. The ML workflow sets that map in motion through iterative D·A·M co-design across data, algorithm, and machine until their emergent capability meets the requirements of the real world.





Learning Objectives

- Explain the six ML lifecycle stages as coordinated data-algorithm-machine decisions with feedback
- Compare ML workflows with traditional software using drift, nondeterminism, and operational feedback
- Analyze how problem-definition constraints propagate through data, modeling, validation, and deployment
- Calculate late-discovery integration costs using the workflow's constraint propagation principle
- Evaluate accuracy, efficiency, reproducibility, and deployment readiness trade-offs across lifecycle stages
- Design feedback loops that connect monitoring signals to retraining, maintenance, and revised requirements

3.1 ML Lifecycle

CONSIDER an illustrative failure scenario. Day one: “Build a diagnostic model for rural clinics.” Day 90: 95 percent accuracy on the test set. Day 120: 96 percent accuracy after a month of architecture tuning. Day 150: model handed to deployment engineers. Day 151: deployment engineers report the model requires 4 GB of memory. Day 152: someone checks the deployment target—tablets in mobile clinics with 512 MB available. Day 153: five months of work is discarded.

The model's accuracy was excellent. The team's machine learning skills were excellent. *The failure was a workflow failure.* A deployment constraint that should have shaped every decision from day one was discovered only after the work was done. The tablet's memory limit should have propagated backward to the first architecture meeting, constraining which models were even worth considering. Instead, the team optimized each component in isolation (data collection, architecture selection, training), and the integration failure appeared only when the pieces were assembled. A documented diabetic retinopathy (DR) deployment exposed analogous workflow and infrastructure gaps.¹

ML systems combine Data, Algorithm, and Machine under physical constraints that partition deployment into Cloud, Edge, Mobile, and TinyML. The parts and operating environments are now in place. The missing piece is how these components connect into a functioning system.

The **ML workflow** is an engineering framework designed to prevent that failure by making constraints explicit at each development stage and tracing how they propagate across Data, Algorithm, and Machine. It marks the shift from model researcher to systems engineer. A researcher optimizes individual elements: a better architecture, a cleaner dataset, a faster accelerator. A systems engineer orchestrates those elements into production systems that reliably deliver value. The day-153 failure was not a data problem, a modeling problem, or a hardware problem in isolation; it was a missing connection among all three. The workflow supplies the mental map that keeps technical decisions attached to the larger system.

The **machine learning lifecycle** is the orchestration framework for this work: a structured, iterative process² that guides the development, evaluation, and improvement of ML systems (Amershi et al. 2019). The formal definition emphasizes continuous management rather than a one-time release.

Here, *lifecycle* describes the stages themselves and *workflow* describes the engineering discipline of orchestrating them; the lifecycle is *what gets traversed*, the workflow is *how the traversal is managed*. This distinction requires **systems thinking**: analyzing how a system's parts interrelate rather than treating them in isolation. The patterns formalized in section 3.9 and illustrated through the detailed case study explain why ML systems require integrated engineering approaches rather than sequential component optimization.

Orchestrating machine learning requires managing data evolution in lockstep with model updates. In figure 3.1, the horizontal layout separates the top data pipeline from the bottom model development pipeline, with curved return vectors illustrating how operational insights and data defects cycle back into earlier stages.

1 | Lab-to-deployment gap:

Beede et al. (2020) documented this empirically for a deep-learning diabetic retinopathy screening system deployed in Thai clinics. Although the system had specialist-level accuracy in earlier validation, 393 of 1,838 submitted images (21 percent) failed its image-quality requirements, often because clinic lighting, camera maintenance, or dilation practices did not match the system's assumptions. Those rejections added work for nurses, forced retakes or referrals, and exposed workflow and infrastructure barriers rather than simply a model-accuracy problem.

2 | CRISP-DM (cross-industry standard process for data mining):

CRISP-DM codified data-intensive system development as six interconnected, iterative phases rather than a linear waterfall (Chapman et al. 2000); its core design principle (feedback loops between all phases) directly informs the modern ML lifecycle's structure. Boehm's software-engineering economics work discussed the higher cost of late fixes (Boehm 1981). The workflow model later uses $2^{N_{\text{stage}}-1}$, where N_{stage} is the lifecycle stage index, as an illustrative sensitivity scenario rather than an empirical cost law.

Definition 3.1: Machine learning lifecycle

Machine learning lifecycle is the iterative engineering process of building, deploying, monitoring, and updating ML systems, where evidence from later stages can feed back to earlier stages because model performance may change after deployment.

1. **Significance:** The lifecycle is a closed loop, not a linear pipeline. Distribution divergence $\mathcal{D}(P_t \| P_0)$ between current and training traffic is an alert signal that raises the probability of accuracy loss; the exact relationship depends on the model, labels, loss, and deployment distribution. A full retraining run incurs the $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$ compute cost, while partial or incremental updates may cost less. Drift velocity and validation delay therefore turn lifecycle maintenance into a budgeting problem, not just an engineering process.
2. **Distinction:** Unlike traditional software, whose behavior changes when code, configuration, dependencies, or environment changes, ML behavior can also change as the world changes. A deployed model's accuracy may change through distribution shift even when the code, infrastructure, and configuration remain untouched.
3. **Common pitfall:** A frequent misconception is that the lifecycle ends at deployment. In reality, deployment is the beginning of the feedback loop: production monitoring surfaces drift, drift triggers investigation, and outcome evidence determines whether a retrained model should re-enter the deployment stage.

Understanding how these stages interconnect is essential before diving into individual layers, as each technical domain—from storage pipelines (chapter 4) to model training (chapter 8), execution frameworks (chapter 7), and serving infrastructure (chapter 13)—operates as an interdependent system.




Figure 3.1: Dual-Pipeline ML Development: Machine learning engineering decouples into parallel data and model pipelines operating on different timescales. While the model pipeline executes rapid inner-loop training and evaluation, outer-loop feedback from production deployment and data validation continuously reshapes upstream collection, labeling, and curation.

The conceptual stages of the ML lifecycle establish the *what* and *why* of the development process. The operational layer constitutes the *how*: the implementation of this lifecycle through automation, tooling, and infrastructure. Chapter 14 names and develops those practices in detail. This distinction matters: the lifecycle is the conceptual framework; operational infrastructure is the machinery that implements it at scale.

3.1.1 Quantifying the ML lifecycle

Practitioner time allocation makes the lifecycle bottleneck measurable: the stages that consume engineering effort are often not the stages that receive the most attention. Understanding the ML lifecycle conceptually is necessary but insufficient for engineering decisions; quantitative characterization reveals where effort and compute actually go in ML projects, exposing which stages bottleneck development and where optimization investments yield the highest returns.

Practitioner surveys reveal a stark asymmetry between where research attention centers and where engineering hours are spent, with data collection and preparation consuming 79 percent of practitioner time. In figure 3.2, a pie chart breaks down reported primary time sinks, contrasting the dominant data collection and preparation slices against the narrow share occupied by model architecture and algorithm tuning.




Figure 3.2: Data Scientist Most-Time Responses: In CrowdFlower’s 2016 survey, 60 percent of respondents selected cleaning and organizing data as the activity where they spent the most time, with 19 percent selecting data collection. Model-focused responses such as pattern mining, training set construction, and algorithm refinement together represent roughly 16 percent of responses. Source: CrowdFlower, *Data Science Report 2016*.

Beyond time allocation, iteration cycles characterize successful ML projects. Return to figure 3.1 and notice the feedback loops driving these iterations: each arrow represents a path that teams traverse repeatedly. Production ML systems usually require repeated iteration across data, model, and infrastructure stages, where each cycle may revisit multiple stages. Understanding what triggers these iterations guides resource allocation. Data quality issues (missing labels, distribution mismatches, preprocessing errors) are often a major source of rework. Architecture and training choices (model capacity, tunable settings such as learning rate and batch size, training instability) and infrastructure issues (latency violations, resource constraints, integration failures) create additional loops that teams must budget for explicitly.


Napkin Math 3.1: The iteration tax

Problem: A DR screening system for rural clinics must choose between a large ensemble trained on high-resolution fundus images (training time: 1 week, accuracy: 95 percent) and a lightweight model suitable for edge deployment on clinic hardware (training time: 1 hour, accuracy: 90 percent). Which model has higher modeled accuracy after six months under these assumptions?

Math: In six months (~26 weeks), the possible experiment count is:

1. Large model: 26 weeks of calendar time at 1 week per experiment. Each experiment improves accuracy by an assumed constant ~0.15 percentage points.
2. Small model: 26 weeks of calendar time at 168 h/week gives 4,368 possible experiments at 1 hour each. Machine time is not the binding constraint at that cycle length, so this scenario counts only 100 of them as *effective* experiments, the number a team can realistically design, run, and learn from in six months. The remaining capacity sits idle for want of hypotheses. Even with a smaller assumed gain per iteration, more useful iterations can produce a larger cumulative gain.

Result: If each iteration improves accuracy by 0.1 percentage points on average, the small model starts at 90 percent and reaches 100 percent after 100 effective iterations, before applying the 99 percent ceiling. It therefore renders as 99 percent. The large model starts at 95 percent and reaches 98.9 percent after 26 slower iterations. In practice, the small model’s rapid iteration enables discovering better architectures, label-preserving data augmentations, and tunable training settings.



Under the stated assumptions, the faster-cycle model finishes slightly ahead after 26 weeks.

Systems insight: *Iteration velocity is a feature.* When candidate systems have comparable attainable quality and useful experiments have similar expected value, shorter cycles expand how much of the design space a team can test. Speed alone cannot guarantee that a smaller model will outperform a larger one. For our DR screening scenario, the lightweight model's rapid iteration cycle enables the team to experiment with label-preserving input transformations, preprocessing pipelines, and architecture variations far more quickly.

These proportions show why data engineering capabilities deserve explicit planning rather than being treated as setup work. They also explain why Part I concludes with chapter 4: data work is a major source of effort, iteration, and project risk. Understanding the data pipeline first provides leverage before the modeling, training, and optimization techniques that follow.

Late discoveries can be costly,³ as formalized later by the constraint propagation principle (section 3.9.1). Violations discovered late may require corrections across multiple preceding stages. This risk motivates explicit **stage interface contracts**: validating outputs at each stage transition catches violations early, while correction costs remain manageable. Section 3.2.2 formalizes these contracts once the six stages themselves have been introduced.

This opportunity cost of slow iteration creates the **iteration tax**. A quick calculation makes the bottleneck concrete.

The iteration tax makes a broader point: *ML workflows are not slow versions of traditional software lifecycles.* They are structurally different, and the differences show up in *where* time is spent, *how* feedback loops operate, and *how* late discoveries can expand rework.

3.1.2 ML vs. traditional software

Traditional and ML systems can both use iterative lifecycles and exhibit nondeterministic behavior.⁴ The relevant distinction is how application behavior is specified. Rule-based software encodes explicit logic, while ML systems learn statistical mappings from data (section 1.6). This adds data, evaluation, and monitoring concerns to established software-engineering practices.

Machine learning systems therefore require additional workflow controls. Consider financial transaction processing: a deterministic authorization rule follows explicit logic and can execute quickly once its inputs are available. An ML-based fraud detector adds learned scoring, feature retrieval, and statistical decision logic. The model stage may take milliseconds, but end-to-end latency also depends on storage, networking, and surrounding services. This shift from explicit programming to learned behavior reshapes the development lifecycle, altering how teams establish reliability and robustness.

These differences alter how lifecycle stages interact. While conventional software engineering relies on production feedback to refine code and requirements, ML systems introduce feedback loops where operational data reshapes training distributions, detected drift prompts investigation, and runtime errors expose dataset blind spots. Table 3.1 contrasts traditional software and ML engineering across six core lifecycle dimensions, highlighting how data evolution redefines testing, deployment, and ongoing maintenance.⁵

Table 3.1: Traditional Software vs. ML Development Lifecycles: Six dimensions where ML development diverges from traditional software engineering. Both use production feedback, while ML systems add loops in which deployment insights reshape data collection, monitoring drives model updates, and production experience informs model design.

Aspect	Traditional Software Lifecycles	Machine Learning Lifecycles
Problem Definition	Precise functional specifications are defined upfront.	Performance-driven objectives evolve as the problem space is explored.
Development Process	Iterative code, configuration, and interface development.	Iterative experimentation with data, features, and models.
Testing and Validation	Many functional tests have exact expected results; performance and reliability remain quantitative.	Statistical validation and metrics that involve uncertainty.
Deployment	Behavior remains static until explicitly updated.	Performance may change over time due to shifts in data distributions.

³ | **Late correction costs:** Boehm's *Software Engineering Economics* (Boehm 1981) discussed how defects found late can cost more to fix than those caught during requirements. In ML systems, late-discovered constraints may require retraining, data-pipeline changes, and renewed validation. The doubling rule is an illustrative sensitivity scenario, not an empirical cost law.

⁴ | **Waterfall model:** A plan-driven lifecycle is often depicted as a sequence of requirements, implementation, and testing, but Royce (1970) warned that a purely sequential implementation was risky and described iteration between successive phases. ML systems add data and model feedback that must be managed explicitly. This chapter uses an illustrative sensitivity model for ML lifecycle stages.

⁵ | **Data versioning:** Unlike code, which changes through discrete, auditable commits, data can drift gradually (distribution shift), suddenly (schema migration), or subtly (label quality degradation). Plain Git is impractical for many multi-terabyte datasets, so teams commonly version external-data manifests or pointers using tools such as DVC and Git LFS. Without data versioning, teams cannot reliably reproduce a prior training run or determine whether an accuracy regression stems from a code change or a data change.

Aspect	Traditional Software Lifecycles	Machine Learning Lifecycles
Maintenance	Maintenance involves modifying code to address bugs or add features.	Continuous monitoring, updating data pipelines, retraining models, and adapting to new data distributions.
Feedback Loops	Production feedback commonly changes code, configuration, and requirements.	Insights from deployment and monitoring often refine earlier stages like data preparation and model design.

 Checkpoint 3.1: ML vs. traditional software

- ML systems are not traditional software with a model attached. Check the differences that force a separate workflow:
- ☐ Failure modes: Can you distinguish silent degradation from explicit failures, recognizing that either can occur in ML and conventional systems?
 - ☐ Logic source: Can you explain why changing training data can change model behavior even when the code remains fixed?
 - ☐ Iteration: Can you explain why ML requires monitoring and, when evidence warrants it, model or data updates?

3.1.3 Data locality at scale

Large ML data pipelines can stress the locality optimizations used by modern operating systems. OS kernels exploit spatial and temporal locality,⁶ the tendency for a program that reads byte *X* to read nearby bytes soon and to reuse recently accessed memory.

When randomized sample order maps to scattered physical reads, shuffling can reduce the effectiveness of file-system buffers and virtual-memory prefetchers. The penalty depends on how logical examples map to physical storage. Contiguous reads from shuffled shards may retain useful locality, whereas small reads across many shards expose storage and network latency. Layout-aware batching, caching, and explicit prefetching can restore locality and reduce these stalls. The relevant systems variable is the physical access pattern presented to the memory hierarchy, not randomization alone.

That distinction determines where to optimize. If the loader already issues large sequential reads, more caching may add little. If it fans out small reads, repacking records or increasing read granularity may matter more than adding compute. Profiling should therefore examine storage request size, queue depth, cache hit rate, and accelerator idle time together. The data-engineering and hardware-aware optimizations examined in the following Parts address these costs. That diagnosis should trace the complete access path from storage to accelerator.

3.2 Lifecycle Stages

The rural-clinic failure shows why ML projects need an explicit six-stage framework: deployment constraints surfaced only after data, model, and evaluation decisions had hardened around the wrong target. Conventional and ML lifecycles both use iteration, while ML systems add feedback from deployment into data, training, and evaluation. The six-stage framework captures this loop.

The complete machine learning lifecycle distills into six core stages laid out across two functional rows. In figure 3.3, trace the top row from problem formulation through model evaluation, and follow the bottom row as deployment feeds monitoring signals back into upstream data collection.

To make these stages concrete, consider how they apply to MobileNetV2, a mobile vision model whose small footprint makes deployment constraints visible. For MobileNetV2, Problem Definition establishes tight constraints: about 14 MB model size, about 600 MFLOP, and real-time inference on mobile-class hardware. Those constraints shape the rest of the workflow immediately. Data Collection must account for on-device preprocessing limitations, and Model Development must choose an architecture whose operations fit the budget. MobileNetV2 does this through depthwise separable convolutions,⁷ not merely by reducing parameter count. Evaluation validates both accuracy

⁶ | **Locality of reference:** Denning (1968) formalized the working-set principle for virtual memory. Randomly ordered examples can reduce locality when the logical order maps to scattered physical reads, but the effect depends on storage layout, batching, caching, and prefetching.

⁷ | **Depthwise separable convolutions:** Replacing a standard convolution with this cheaper factorization reduces computation by roughly 8–9× for typical kernel sizes (Sandler et al. 2018), which is what makes the roughly 600 MFLOP inference budget plausible on mobile-class hardware. Chapter 6 covers the architectural mechanism in depth.




Figure 3.3: Simplified Lifecycle with Feedback: The six core stages of ML engineering form a closed loop rather than a linear handoff. A primary feedback path from monitoring back to data collection supports investigation and adaptation as production distributions shift, model accuracy degrades, and operational constraints evolve.

and latency on target devices, Deployment tests whether the model fits the device’s memory and power envelope, and Monitoring tracks performance across diverse device populations. Each stage’s decisions propagate through subsequent stages, and the workflow framework makes these dependencies explicit. A DR screening model optimized for rural clinic deployment faces analogous pressures: limited device memory, strict power budgets, and the need for real-time inference without reliable connectivity. These shared constraints make DR an effective running case study.

Figure 3.3 suggests linear progression, but the feedback loop reveals the true iterative nature of ML development.

Checkpoint 3.2: The workflow cycle

The stage prompts test whether the lifecycle can be traced as a coupled system.

The Stages

- ☐ Problem definition: Have you defined success metrics that actually map to business value?
- ☐ Data: Is your data pipeline reproducible, and have you specified what it must deliver before model development begins?
- ☐ Modeling: Are you iterating fast enough?
- ☐ Deployment: Have you recorded the deployment constraints that must shape upstream data and model decisions?

Figure 3.3 captures this loop, but its deeper weight is quantitative: each stage corresponds to specific terms in the performance equation, and this mapping reveals a workflow-level version of the iron law: decisions made during data collection constrain what is achievable during model development, which in turn determines deployment requirements. The stage-by-stage mapping connects the lifecycle to the iron law of ML systems defined in section 1.7.

The binding constraint differs dramatically across workload archetypes, causing each lifecycle stage to optimize different iron law terms. ResNet-50, DLRM, and keyword spotting (KWS) are useful anchors because each stresses a different part of the system: dense vision training tries to keep accelerators busy, sparse recommendation spends much of its time moving embedding rows, and keyword spotting is constrained first by tiny memory and always-on energy budgets. Table 3.2 shows how three selected workflow stages manifest for these recurring workload archetypes.

Production systems rarely fall neatly into a single archetype. A medical imaging classifier, for instance, may require sustained computation while being trained over large image datasets yet face strict energy and memory constraints when deployed to portable clinic devices. Understanding *how*

the same workflow framework adapts to each archetype, and *how* a single project can span multiple archetypes simultaneously, is essential for making sound engineering decisions.

Table 3.2: Workflow Variations by Lighthouse Model: Three selected lifecycle stages target different iron law terms depending on the workload’s binding constraint. ResNet-50 emphasizes useful computation per unit time, DLRM emphasizes data movement and embedding freshness, and KWS emphasizes energy and memory capacity.

Stage	ResNet-50 vision training	DLRM recommendation	KWS TinyML
Data Eng	Keep image batches available fast enough to sustain high GPU utilization.	Keep online feature-store lookups within the serving latency budget; embedding tables dominate storage and freshness.	Curate short audio clips for an SRAM-constrained device.
Training	Coordinate preprocessing, batching, mixed precision, and model execution to reduce accelerator idle time.	Optimize sparse embedding lookups because memory bandwidth limits throughput.	Search for the smallest model family that still recognizes the trigger phrase reliably.
Deploy	Use large batches when throughput and cost matter more than single-request latency.	Meet strict interactive p99 latency targets while keeping features fresh.	Stay within a strict always-on energy budget.

Each stage of this workflow presents distinct engineering challenges, from curating high-quality datasets to maintaining model performance in production. DR screening earns its place as the case study that threads through every stage by passing three tests: it appears simple on the surface but reveals deep complexity in practice, it spans enough of the deployment spectrum to exercise the workflow framework, and documented studies can be paired with explicit illustrative scenarios.

🔍 Systems Perspective 3.1: The iron law of workflow

The six lifecycle stages are not merely procedural steps. Each can alter several variables in the iron law of ML systems ($T = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}}$), so the mappings below are diagnostic starting points rather than one-to-one assignments:

- Problem definition: Sets the target constraints: accuracy, latency, cost, privacy, and deployment paradigm. These targets determine which terms of the equation are allowed to grow and which must be bounded from the start.
- Data collection and preparation: Shapes dataset composition and the byte volume (D_{vol}) presented to downstream pipelines. Curation can also change the work required to reach a quality target.
- Model development and training: Changes the operation count (O), data reuse, parameter movement, and achievable hardware efficiency (η_{hw}).
- Evaluation and validation: Tests whether model quality and end-to-end performance meet deployment requirements on the target system.
- Deployment and integration: Determines the serving path, including data movement, execution efficiency, and fixed overhead (L_{lat}).
- Monitoring and maintenance: Observes drift, latency, throughput, and cost after launch, then feeds violations back into earlier stages for re-optimization.

When we view the lifecycle stages through the iron law, the equation provides one quantitative lens on latency and cost. Workflow management must also account for accuracy, safety, privacy, and organizational constraints outside the expression.

3.2.1 Case study: DR screening

The documentation behind DR screening spans both sides of the lab-to-field divide: Gulshan et al. (2016) document the large validation study for automated DR detection, while Beede et al. (2020) document what changed when a related deep-learning system was used in Thai clinics. The

problem appears straightforward (classify retinal images as healthy or diseased), but the path from laboratory success toward clinical use illustrates lifecycle complexity. Together, these sources give us a documented path from data collection and validation to workflow integration and infrastructure constraints.

Diabetic retinopathy is a leading cause of preventable blindness.⁸ To appreciate what the model must learn, look closely at figure 3.4: the clinical challenge is detecting characteristic hemorrhages (dark red spots) that indicate disease progression. Limited specialist access in some regions motivates scalable screening programs in which AI assistance may play a role.




Figure 3.4: Retinal Hemorrhages in DR Screening: Side-by-side fundus images illustrate the subtle visual features an ML screening system must detect. Identifying localized microvascular lesions (such as the indicated hemorrhage) requires preserving fine spatial resolution through the vision pipeline and establishing strict input-quality validation to handle lighting and camera variations in clinical field deployments. Source: Google.

Initial research achieved expert-level performance in controlled settings. However, the journey to clinical deployment revealed how technical excellence must integrate with data quality challenges, infrastructure constraints in rural clinics, regulatory requirements, and workflow integration.⁹ For the metrics that recur in this case, sensitivity is the true-positive rate, specificity is the true-negative rate, and AUC is the area under the ROC curve. The same constraint propagation dynamics apply whether the target is medical imaging systems or mobile applications like MobileNetV2, and the lifecycle stages ahead trace these dynamics in concrete detail.

3.2.2 Stage interface specification

Each lifecycle stage operates as a distinct engineering phase with defined inputs, outputs, and quality invariants. Think of these as *API contracts* between teams. A microservice relies on an interface specification to make integration assumptions explicit; a data pipeline similarly relies on schema and distribution contracts to detect incompatibilities before they propagate. Table 3.3 formalizes these contracts, making explicit what each stage must receive and produce. This turns the abstract lifecycle diagram (figure 3.3) into actionable engineering requirements. When a stage’s output fails to meet its contract, the deficiency may propagate into dependent stages.

Table 3.3: Stage Interface Specification: Each lifecycle stage has explicit input requirements, output deliverables, and quality invariants that must hold for the stage to be considered complete. Violations of these contracts create technical debt that compounds through subsequent stages. The deployment paradigm selection in Problem Definition (Cloud, Edge, Mobile, or TinyML from chapter 2) constrains all downstream stages, as a TinyML target imposes different data, model, and monitoring requirements than a Cloud target.

Stage	Input Contract	Output Contract	Quality Invariant
Problem Definition	Business requirements; operational context	Measurable objectives; deployment paradigm selection; resource constraints	All success criteria are quantifiable; target deployment paradigm is explicit
Data Collection & Preparation	Objectives; deployment target; quality requirements	Versioned dataset with schema; preprocessing pipeline; data validation rules	Distribution approximates anticipated production environment; labeling meets accuracy requirements

8 | Diabetic retinopathy (DR): DR is a common complication of diabetes; early detection can prevent vision loss, but specialist access varies substantially across countries. This gap motivates scalable screening programs. In clinics with limited hardware or unreliable connectivity, model size, offline capability, latency, and workflow integration can become binding deployment constraints.

9 | Healthcare AI deployment gap: In the Thai DR deployment studied by Beede et al. (2020), real clinics exposed workflow, image-quality, infrastructure, and human-factors constraints that were absent from controlled evaluation. The case shows why laboratory metrics alone do not establish deployment readiness.

Stage	Input Contract	Output Contract	Quality Invariant
Model Development & Training	Dataset; accuracy targets; resource constraints	Trained model weights; training configuration; experiment logs	Meets accuracy thresholds within computational budget; architecture compatible with deployment target
Evaluation & Validation	Trained model; held-out test data; evaluation criteria	Performance metrics across subgroups; failure mode analysis; validation certificate	No critical subgroup falls below minimum thresholds; confidence calibration meets domain requirements
Deployment & Integration	Validated model; infrastructure requirements; service-level agreement (SLA) targets	Serving endpoint; monitoring instrumentation; rollback procedures	Latency and throughput meet paradigm requirements; integration tests pass
Monitoring & Maintenance	Live system; performance baselines; alert threshold	Drift detection alerts; update-review criteria; incident reports	Alert coverage and detection windows are defined; detected degradation triggers review

This specification reveals why ML projects experience the iteration cycles diagrammed in figure 3.3. When a downstream stage discovers that an upstream contract was violated (for example, evaluation reveals the training data distribution does not match production), the project must iterate back to fix the root cause. Teams that validate contracts at each stage transition catch violations early, when correction costs are lowest. This validation process is best understood as *auditing stage transitions*.



Example 3.1: Auditing stage transitions

Scenario: An engineering team claims to have completed Problem Definition for a medical imaging classifier, attempting to transition to Data Collection without establishing deployment target constraints.

Diagnosis: The output contract lacks deployment paradigm selection and resource constraints (latency and memory targets). Attempting to deploy late in the lifecycle exposes violations (e.g., edge device memory limit < 200 MB), incurring 16× the cost compared to early correction.

Systems lesson: Stage transitions act as control-plane quality gates in the ML lifecycle. Enforcing strict output contract validation before proceeding downstream prevents expensive upstream rework and ensures data collection aligns with target deployment hardware.

The DR case study and Stage Interface Specification provide the concrete context and formal contracts that ground each lifecycle stage. The first stage, Problem Definition, records the constraints that subsequent stages must satisfy.

3.3 Problem Definition

Problem definition in ML begins with sentences that look deceptively simple. A product manager writes: “Build a model that detects diabetic retinopathy.” That single sentence conceals a dozen engineering decisions: sensitivity thresholds for patient safety, hardware capabilities in rural clinics, latency budgets that keep clinicians engaged, and regulatory frameworks governing approval. Some conventional requirements translate directly into implementation rules. In ML systems, defining *what* the system should do is inseparable from defining *how* it will learn to do it and the physical constraints under which it must operate. This first stage, the leftmost box in figure 3.3, lays the foundation for all subsequent phases in the ML lifecycle.

The DR screening case makes this concrete. What appears to be a straightforward classification task (detect disease in retinal photographs) actually requires balancing five competing constraints: diagnostic accuracy (patient safety), computational efficiency (rural clinic hardware), workflow integration (clinical adoption), regulatory compliance (FDA approval), and cost-effectiveness (sustainable deployment in resource-limited settings). Each constraint tightens the feasible design space for the others: pursuing higher accuracy through larger models conflicts with the hardware budget; achieving regulatory compliance demands annotation protocols that increase data collection costs. ML adds learned statistical behavior to a multi-constraint optimization problem that also appears in conventional systems.

3.3.1 Constraint layers

The DR example reveals that ML problem definitions are not single requirements but stacks of interacting **constraint layers**. Accuracy constraints (>90 percent sensitivity, >80 percent specificity across diverse populations and equipment) sit on top of infrastructure constraints (edge devices with limited compute, intermittent connectivity, inference within clinical workflow timeframes) which sit on top of regulatory constraints (FDA validation, audit trails, privacy compliance). Each layer narrows the feasible design space for the layers above it.

Privacy compliance in an ML system carries a distinctive operational weight that a generic data-handling rule does not capture. In a traditional database, deleting a patient's record is a DELETE statement. In an ML system, model weights may encode statistical patterns learned from that record: where law or a regulator requires removing that influence, compliance may require machine unlearning (an active research area with incomplete guarantees) or retraining on the remaining data. The obligation and remedy depend on the governing rule and system; at DR-system scale, either can make privacy compliance a recurring compute-budget item and training provenance an architectural constraint.

This layered structure generalizes beyond healthcare. Any ML problem definition must address at least three constraint layers: statistical (what accuracy, across which subpopulations), physical (what hardware, under what latency and memory budgets), and operational (what regulatory, organizational, or workflow requirements apply). The constraint propagation principle (section 3.9.1) explains why: a constraint that exists but remains unspecified does not disappear—it may surface later after dependent decisions have hardened.

The specific constraints for the DR system did not emerge from technical analysis alone. They required systematic collaboration between engineers, ophthalmologists, and clinic administrators to translate clinical needs into measurable engineering requirements. Key decisions (balancing model complexity with hardware limitations, ensuring interpretability for healthcare providers, and accounting for patient privacy) emerged from this cross-disciplinary process. Without domain expertise, the engineering team might have optimized for aggregate accuracy while missing the sensitivity threshold that determines clinical safety.

War Story 3.1: When the label was the bias (2018)

Context: In 2014, Amazon assembled an engineering team at its Edinburgh hub to build an internal machine-learning system for ranking technical job applicants, training it on ten years of resumes submitted to Amazon (Dastin 2018).

Mechanism: Because prior hiring was male-dominated, the model learned that male-coded language correlated with hiring success, systematically penalizing terms like “women’s” and downgrading graduates of all-women’s colleges.

Impact: The model reproduced historical hiring bias, rendering automated resume screening unsuitable for production recruitment.

Fix: Amazon abandoned the experimental tool after attempts to remove explicit gender signals could not guarantee that it would not learn other discriminatory proxies.

Systems lesson: Problem definition must specify fairness, auditability, and rejection criteria before data collection and training. A workflow trained on biased labels does not learn the operational goal; it learns to reproduce the labels.

3.3.2 Problem definitions evolve

ML problem definitions may evolve as a system scales. Suppose a DR program grows from a handful of clinics with consistent imaging setups to hundreds with varying equipment, staff expertise, and patient demographics.¹⁰ Such growth may require stratified accuracy targets, support for heterogeneous hardware, and revised reporting requirements.

This evolution is not a sign of poor initial planning—it is inherent to ML systems. Scaling exposes edge cases invisible at pilot scale, and production data reveals distributional properties that no training set fully captures. The problem definition must accommodate this reality by specifying

10 | **Demographic drift:**

Population changes can expose performance gaps that aggregate evaluation obscures. Facial-analysis benchmarks found large demographic error-rate disparities (Buolamwini and Gebru 2018), while a health-care risk-scoring study found racially biased allocation from a proxy label despite similar need (Obermeyer et al. 2019). These examples illustrate why a growing DR program should evaluate subgroup performance rather than claiming that such a rollout occurred.

both current targets and the mechanisms for revising them: which metrics trigger re-evaluation, who approves revised thresholds, and how changes propagate to downstream stages.

3.4 Data Collection

With objectives defined and constraints layered, the next practical task is identifying the data that can teach the model to meet these objectives. The constraints, metrics, and deployment targets from problem definition exist only on paper until a team acquires the data that will teach the model to satisfy them. This transition from defining goals to data collection marks a critical juncture where many projects fail. As the survey in section 3.1.1 established, practitioners often report data-related activities as their primary time sink. In iron law terms, this stage primarily determines dataset size and composition (D), along with the byte volume (D_{vol}) that downstream stages must move. The deployment constraints established during problem definition now become data requirements: if the model must run on edge devices, the data pipeline must produce inputs compatible with edge preprocessing. If the model must achieve 90 percent sensitivity across diverse populations, the data must include sufficient examples from each population.

Data collection and preparation is not merely a preliminary step but often a major engineering activity. Chapter 4 addresses data engineering as its core focus. For DR screening, the challenge is substantial: the data must be statistically diverse enough to train a model that generalizes across populations, operationally feasible to collect in resource-limited clinics, and annotated with enough clinical rigor to satisfy regulatory scrutiny.

Problem definition decisions shape data requirements in the DR example. The multi-dimensional success criteria established (accuracy across diverse populations, hardware efficiency, and regulatory compliance) demand a data collection strategy that goes beyond typical computer vision datasets. Not all data contributes equally to learning, either—chapter 9 shows how selection can retain comparable performance with less compute when a dataset contains redundancy, provided the selected subset preserves the evidence the task requires.

The DR system requires on the order of 10^5 retinal fundus photographs, each reviewed by multiple expert ophthalmologists. Expert consensus addresses the inherent subjectivity in medical diagnosis (two ophthalmologists may disagree on borderline cases) while establishing ground truth labels that can withstand regulatory scrutiny. The annotation process must capture clinically relevant features like microaneurysms, hemorrhages, and hard exudates across the full spectrum of disease severity.

Raw high-resolution retinal scans can reach tens of megabytes per image before compression, creating substantial infrastructure challenges. A clinic processing dozens of patients per day can produce gigabytes to tens of gigabytes of imaging data per week. This data volume can occupy much of a limited uplink during clinic hours, making local processing one option alongside compression, batching, scheduling, or increased link capacity.



Napkin Math 3.2: Bandwidth vs. compute

Problem: A rural clinic captures retinal images for DR screening. How much of the assumed same-shift uplink window would raw uploads occupy, and what alternatives reduce that occupancy?

Math:

1. Daily data: At 150 patients/day, 10 photos/patient, and 5 MB/photo, the clinic produces 7.5 GB/day.
2. Upload time: The uplink is 2 Mb/s, or 0.25 MB/s. Dividing 7,500 MB by that rate gives 30,000 s, approximately 8.3 h.
3. **Constraint:** If the clinic requires same-shift upload during its 8 h operating window, this transfer alone occupies the link for 104.2 percent of that window and may contend with other traffic.

Systems insight: Local inference with uploaded detection summaries (10 KB/patient) reduces bandwidth usage by 5,000×. Overnight batching, compression, scheduling, or a faster link are alternatives not compared by this calculation.

~5000×

Raw

Summary

log scale

Edge summaries beat raw medical-image uploads when bandwidth dominates.

3.4.1 Lab-to-field data gap

Laboratory data and production data inhabit different worlds. This **lab-to-field gap** appears when DR screening deploys to rural clinics across Thailand and India: images arrive from diverse camera equipment operated by staff with varying expertise, often under suboptimal lighting with inconsistent patient positioning. A model trained on high-quality research images from standardized fundus cameras may fail on blurry, poorly-lit images from older equipment—not because the algorithm is wrong, but because the data distribution has shifted beyond the training envelope.

As the Bandwidth vs. Compute exercise quantified, this data volume makes same-shift raw upload difficult under the scenario assumptions. The scenario therefore evaluates edge deployment using specialized hardware such as NVIDIA Jetson.¹¹ Local preprocessing reduces bandwidth requirements by orders of magnitude but demands correspondingly more local computation, forcing a trade-off: simpler models that run on constrained hardware, or more powerful edge devices that increase per-clinic costs.

The bandwidth constraint makes infrastructure a data-collection decision rather than a late implementation detail. The scenario architecture combines edge devices for local inference and preprocessing, clinic aggregation servers for data management and buffering, and cloud training infrastructure for periodic model updates. It targets end-to-end latency under 100 milliseconds and operation without connectivity-induced delays.

Privacy constraints impose a similar architectural decision. Patient privacy regulations may motivate local retention or distributed training, but they generally specify protections and permitted uses rather than one required training topology. Workflows that keep raw clinic data local while sharing only approved updates or summaries add complexity to both data collection and model training infrastructure.

3.4.2 Distributed data infrastructure

If a deployment grows from a handful of clinics to hundreds, data infrastructure must scale accordingly. Each retinal image travels through multiple stages: clinic cameras capture the image, local systems provide initial storage and processing, quality validation checks ensure usability, secure transmission moves data to central systems, and finally, integration with training datasets completes the pipeline. The infrastructure decisions at each stage are shaped by the deployment constraints established during problem definition.

Storage tiers are another place where the data pipeline either preserves or erodes downstream iteration velocity. Different data access patterns demand different storage solutions, so teams typically implement **tiered storage architectures**,¹² each calibrated to access frequency and performance requirements.

Hot storage uses high-throughput NVMe SSDs for data currently used in training loops. Warm storage uses S3-compatible object storage for recent datasets and active validation sets. Cold storage uses low-cost archival systems, such as AWS Glacier, for historical data required for regulatory audit trails but rarely accessed.

In practice, the boundary between tiers is dynamic: a dataset migrates from warm to hot when selected for the next training run, and from hot to cold when the model it trained is superseded. Automated lifecycle policies manage these transitions, promoting data based on training schedules and demoting it based on access recency—a pattern that chapter 4 explores in detail.

Rural clinic deployments face severe connectivity constraints that force a choice between transmission strategies. Clinics with reliable broadband can stream images in near-real-time for centralized processing, but clinics with intermittent satellite links, common in remote regions of India and sub-Saharan Africa, require **store-and-forward architectures** that batch images during connectivity windows and reconcile results asynchronously. The choice propagates through the entire stack: store-and-forward clinics need larger local storage buffers, more robust local inference capabilities, and conflict-resolution logic when locally generated predictions differ from later cloud-based analysis.

Infrastructure scalability poses a harder challenge than raw capacity. As the system grows from a handful of pilot clinics to hundreds of production sites, data heterogeneity grows alongside data volume: each clinic's camera model, lighting environment, and operator habits can produce a different image distribution. The infrastructure must handle increasing throughput while also

¹¹ | **NVIDIA Jetson:** NVIDIA's Jetson family spans a wide SKU spectrum, from Jetson Orin Nano (7–15 W) through Jetson Orin NX (10–25 W) to Jetson AGX Orin (15–60 W). This scenario uses an Orin Nano-class device (4–8 GB shared LPDDR5, 7–15 W power envelope) as a concrete edge target. Its memory and power budgets constrain model complexity, while selecting a larger device changes both the feasible model and the per-clinic cost.

¹² | **Tiered storage:** Places data on different storage media based on access frequency and performance requirements. The storage price gap is roughly 4.3× in this example: Non-Volatile Memory Express (NVMe) SSDs deliver 500,000+ input/output operations per second (IOPS) at ~\$0.10/GB/month, while object storage costs ~\$0.023/GB/month but with 100–200 ms latency. For ML training loops requiring sustained sequential reads at 1–10 GB/s, choosing the wrong tier converts a compute-bound training pipeline into an I/O-bound one, directly inflating the iron law's data term (D_{vol}/BW).

3.5 Model Development

The DR team has 128,000 labeled retinal images, a validated preprocessing pipeline, and a target: more than 90 percent sensitivity on edge hardware with less than 50 ms inference latency. The question is no longer *what data* to collect but *what model* to build—and that question has no answer independent of the deployment constraints already established. In iron law terms, this stage defines the Operations (*O*) term: architectural choices set the computational floor that hardware must sustain. The challenges extend well beyond selecting algorithms and tuning hyperparameters.¹⁴ Chapter 8 covers the training methodologies, infrastructure requirements, and distributed training strategies in detail. In high-stakes domains like healthcare, every design decision affects clinical outcomes, so technical performance and operational constraints must be integrated from the start.

The DR system faces a sharp training challenge: achieve expert-level diagnostic accuracy with finite labeled-data and optimization budgets. **Transfer learning**¹⁵ addresses this constraint by adapting a model pretrained on a source task to a target task. The ImageNet Large Scale Visual Recognition Challenge (ILSVRC) 2012 training subset contains 1.3 million labeled images (Russakovsky et al. 2015). Reusing learned representations can improve target-task optimization and generalization relative to random initialization, with the benefit depending on source-target similarity, which layers are transferred, and transfer-related optimization effects (Yosinski et al. 2014).

In the controlled validation study reported by Gulshan et al. (2016), transfer learning and a labeled dataset of 128,000 images produced an AUC¹⁶ of 0.99, with sensitivity of 97.5 percent and specificity of 93.4 percent at one operating point. The result demonstrates the potential of large-scale pretraining with domain-specific fine-tuning under the study conditions; it does not by itself establish deployment performance. The training strategy uses error gradients to adjust model weights; chapter 5 establishes that mechanism before chapter 8 develops the training systems around it.

Achieving high accuracy is only the first challenge. Edge deployment constraints impose strict efficiency requirements: models may need to fit within tens to hundreds of megabytes, complete inference in tens of milliseconds, and operate within tight memory budgets.

From a workflow perspective, accuracy gains must always be weighed against deployment feasibility. **Ensemble learning**¹⁷ illustrates this trade-off: combining predictions from multiple models often yields better performance than any individual model, but adds inference work and memory usage. **Bagging** trains multiple models on different data subsets, **boosting** sequentially trains models to correct previous errors, and **stacking** uses a meta-model to combine base model predictions. These methods generate diversity in different ways, but every constituent model adds serving work that must fit the deployment budget. The Netflix Prize illustrates how ensemble complexity can impede production deployment.¹⁸

Initial research models are often much larger (sometimes multiple gigabytes when using ensembles) and therefore violate deployment constraints, requiring systematic optimization to reach a deployable form factor while preserving clinical utility. These constraints drive systematic model compression and optimization rather than isolated accuracy tuning. Later model-compression techniques reduce computation and memory footprint, but each reduction must be revalidated against clinical accuracy. Chapter 10 details those techniques. The development process requires continuous iteration between accuracy optimization and efficiency optimization: model capacity, preprocessing choices, and execution cost all affect both dimensions simultaneously.

3.5.1 Reproducible system artifacts

The accuracy-efficiency balancing act produces more than trained weights alone. A common failure mode is treating the trained model weights as the sole output of this stage. In a mature ML workflow, the deliverable is a **reproducible system artifact** with four components:

- **Model Weights:** The learned parameters.
- **Inference Code:** The exact code used to run the model, including preprocessing logic.
- **Environment Specification:** The complete dependency graph (for example, Docker container, `requirements.txt`, CUDA drivers) required to execute the code.
- **Configuration:** Hyperparameters and runtime settings.

¹⁴ | **Hyperparameter:** Architectural and optimizer choices (for example, learning rate and network depth) affect the computational work of training. A naive Cartesian grid that trains every combination independently incurs a multiplicative search cost: 5 hyperparameters with 4 values each produces 1,024 (4^5) configurations. Early stopping, multifidelity methods, and adaptive search can avoid completing every run, which is why the grid represents a baseline rather than an unavoidable cost.

¹⁵ | **Transfer learning:** For the same architecture, transfer learning changes how parameters are initialized and optimized, not the operations required at inference.

¹⁶ | **AUC (area under the ROC curve):** Measures the area under the receiver operating characteristic (ROC) curve plotting true positive rate vs. false positive rate across all classification thresholds, ranging from 0 to 1; 0.5 represents random discrimination, and values below 0.5 are possible. Unlike accuracy, AUC is threshold-independent and insensitive to class prevalence, making it a common metric for medical screening systems. The systems consequence: a model with 0.99 AUC can still produce unacceptable sensitivity at the specific operating threshold chosen for deployment, so AUC alone cannot validate deployment readiness.

¹⁷ | **Ensemble learning:** Combines predictions from multiple models (bagging, boosting, stacking). Inference compute and model storage generally sum across constituents. Serial execution can increase latency, while parallel execution trades latency for additional hardware and coordination; no fixed proportional latency multiplier follows from ensemble size.

¹⁸ | **Competition-production gap:** The Netflix Prize winner used a complex ensemble, but Netflix did not deploy the winning approach because the additional accuracy did not justify the engineering effort (Johnston 2012).

Without bundling the environment with the model, dependency mismatches can create dangerous failures. Some incompatibilities, including unsupported CUDA combinations or missing libraries, fail loudly and immediately. Other differences in compatible kernels, linear algebra libraries, or image-resizing routines such as OpenCV and PIL may alter floating-point results or pixel interpolation without crashing the serving process. These changes can shift model outputs and affect accuracy. Environment reproducibility is therefore necessary not only for successful execution but also for validating equivalent inference behavior. A model that achieves 99 percent accuracy in development must be reevaluated if its production environment changes, even when inference completes without an exception.

¹⁹ | **Medical AI performance metrics:** Medical AI evaluation emphasizes sensitivity (true positive rate) and specificity (true negative rate) alongside aggregate accuracy. This DR scenario sets a >90 percent sensitivity target because missed referable disease can delay care and contribute to avoidable vision loss; actual acceptance criteria are product- and regulator-specific. The subtler systems trap is positive predictive value (PPV): even a high-accuracy model can have low PPV in a low-prevalence population. This prevalence dependence can require different operating thresholds across deployment sites, a constraint invisible in standard ML evaluation.

²⁰ | **Model compression pipeline:** Bridging the gap between research accuracy and edge deployment requires an iterative “compress-validate-adjust” loop. Each compression step can reduce model size or execution cost, but it can also silently degrade sensitivity below the clinical threshold. Finding a model that fits in device memory while preserving clinical sensitivity typically requires multiple iterations because only the full validation suite reveals whether the smaller model still satisfies the problem definition.

²¹ | **Ablation studies:** Named for surgical tissue removal, ablation studies systematically disable individual components to isolate their contribution to performance. The rigor matters because a 0.5 percent accuracy difference can determine whether the DR model meets its clinical sensitivity threshold. Without ablation, a team cannot distinguish a genuine architectural improvement from noise introduced by a different random seed, wasting iteration cycles on phantom gains.

3.5.2 Accuracy vs. efficiency

Medical applications demand specific performance metrics¹⁹ that differ from the standard classification outputs and losses chapter 5 introduces. A DR system requires high sensitivity (to limit missed referable disease) and high specificity (to avoid overwhelming referral systems). These metrics must be maintained across diverse patient populations and image quality conditions.

Optimizing for clinical performance alone is not enough. In this scenario, the model artifact must remain below 500 MB so that it fits comfortably within the edge target’s 4–8 GB of shared system memory, while the device operates inside a sub-20 watt power envelope and meets the clinical latency budget. Improvements in one dimension often come at the cost of others: the Operations (O) term, the model byte footprint that contributes to D_{vol} , and fixed serving overhead (L_{lat}) can pull in different directions. Chapter 6 explores model capacity, while chapter 2 discusses deployment feasibility, and the inherent tension between them drives architectural decisions. Systematic compression and revalidation²⁰ can bridge the gap, meeting deployment requirements while aiming to preserve clinical utility.

The ensemble trade-off illustrates a broader pattern: choosing an ensemble of lightweight models over a single large model reduces per-model complexity (enabling edge deployment) but increases pipeline complexity (requiring orchestration logic and multi-model monitoring). Every architectural decision creates this kind of downstream ripple.

3.5.3 Constraint-driven development

Real-world constraints shape model development from initial exploration through final optimization, demanding systematic experimentation. Development begins when data scientists collaborate with domain experts (ophthalmologists in the DR case) to identify subtle lesions and image-quality requirements that matter clinically. Without that domain knowledge, a model architect might choose a resolution or receptive field, the input region each internal feature can see, that discards relevant detail before the network can use it. This interdisciplinary approach helps model architectures preserve clinically relevant evidence while respecting the computational constraints identified during data collection.

Computational constraints profoundly shape experimental approaches. Multiple model variants, hyperparameter sweeps, and preprocessing approaches can make exhaustive experimentation expensive. This economic reality drives investments in better job scheduling, caching of intermediate results, early stopping, and automated resource optimization. Systematic hyperparameter optimization and disciplined experiment design can reduce computation relative to exhaustive search.

The inherent uncertainty of ML outcomes demands scientific methodology: controlled variables through fixed random seeds and environment versions, systematic **ablation studies**²¹ to isolate component contributions, confounding factor analysis to separate architecture effects from optimization effects, and statistical significance testing across multiple training runs using paired offline tests rather than the A/B testing²² reserved for comparing models on live production traffic. Without this rigor, teams cannot distinguish genuine performance improvements from statistical noise—a distinction that becomes critical when a 0.5 percent accuracy difference determines whether a model meets the clinical sensitivity threshold.

At every development milestone, teams validate models against the deployment constraints identified in earlier lifecycle stages. Each architectural innovation must be evaluated for accuracy improvements and compatibility with edge device limitations and clinical workflow requirements.

This dual validation approach ensures that development efforts align with deployment goals rather than optimizing for laboratory conditions that do not translate to real-world performance.

3.5.4 Prototype to production

A team of three data scientists may manage experiments with spreadsheets and shared notebooks, while a team of 30 is more likely to need shared workflow infrastructure. As projects evolve from prototype to production, complexity grows across multiple dimensions simultaneously: larger datasets, more sophisticated models, concurrent experiments, and distributed training infrastructure. Informal coordination can become a bottleneck at production scale as teams share data splits, resolve experiment conflicts, and reconcile notebook versions. Figure 3.6 illustrates one assumed contrast between manual coordination and a shared workflow platform. The axes are relative units intended to show shape, not absolute throughput.




Figure 3.6: Illustrative Workflow Automation Dynamics: Conceptual scaling curves contrast the experimentation velocity of manual workflows against shared workflow platforms as team size grows. Unstructured ad hoc scripts eventually hit a “coordination tax” from merge conflicts and unversioned artifacts, whereas shared platforms can reduce duplicated work through reusable components and automated lineage. The curves illustrate assumptions rather than empirical scaling laws.

The curves in figure 3.6 encode assumed functions rather than measured team-scaling laws. The manual-workflow curve represents coordination costs from sharing data splits, resolving experiment conflicts, and reconciling notebook versions. The platform curve represents benefits from reusable preprocessing components, versioned experiment tracking, and automated scheduling. Actual throughput need not saturate or grow super-linearly; it depends on workload parallelism, platform overhead, team structure, and experiment quality. Figure 3.6 therefore illustrates why shared infrastructure can reduce duplicated work and manual handoffs without predicting a universal return for every team.

3.5.5 Reproducibility and technical debt

Shared infrastructure can accelerate experimentation—but rapid iteration creates a hidden liability. If experiments are not reproducible, the team cannot reliably distinguish genuine improvements from noise, and the codebase accumulates technical debt²³ that compounds with every unreproducible result.

Reproducing ML results requires tracking data versions, random seeds, hardware configurations, and library versions alongside program logic. Hardware and nondeterministic operations can alter training results, so teams need lineage to distinguish architectural changes from run-to-run variation. Systematic **experiment tracking** records unique run identifiers and artifact versions. Systems such as MLflow and Weights & Biases link data versions, code commits, hyperparameters, and resulting models so teams can reconstruct differences between runs.

The cost of neglecting **reproducibility** is economic, not just scientific. Teams that cannot reproduce a result waste cycles rerunning experiments that may not converge to the same outcome.

²² | **A/B testing in ML:** Compares a new model (B) against the production baseline (A) on randomly assigned live traffic to estimate the model’s causal effect. Required sample size depends on the baseline event rate, effect size, variance, allocation, significance level, and statistical power; no universal interaction count or duration follows from a 0.5 percentage-point change alone.

²³ | **ML artifact interdependence:** ML artifacts are deeply interdependent. A model’s measured accuracy is tied to the data version, preprocessing pipeline, and hyperparameter configuration used to produce it. Without lineage tracking, the team may not be able to determine whether a regression stems from code, data, or run-to-run variation, forcing expensive reruns of experiments whose provenance is lost.

Reproducibility infrastructure such as versioned environments, controlled pipelines, and automated checkpointing reduces redundant computation and supports more confident architectural decisions.

Reproducible, optimized models are necessary but not sufficient. A model that achieves expert-level accuracy on curated research data may still fail in production. The next stage subjects these trained artifacts to systematic testing against the conditions they will actually encounter.

3.6 Evaluation and Validation

In this illustrative scenario, the DR team’s model achieves an AUC of 0.99 on the curated research dataset. Then they test it on images from a rural clinic in Chiang Mai where a technician with two weeks of training operates a five-year-old fundus camera. Sensitivity drops to 78 percent. The code has not failed or thrown an exception; the model has simply never seen images this blurry, this poorly lit, or this inconsistently framed. Laboratory success does not guarantee production value, and the gap between the two is where many ML projects fail. Before deployment, trained models must undergo rigorous evaluation and validation to confirm they meet performance requirements across the diverse conditions encountered in production. This stage bridges model development and deployment, transforming experimental artifacts into production-ready systems through systematic testing against predefined metrics, edge cases, and real-world scenarios.



Definition 3.2: Model validation


Model validation is an evidence gate for deciding whether a trained model is suitable to deploy for a specified use. It tests the model against deployment constraints such as latency targets, subgroup performance targets, cost budgets, and robustness under distribution shift, but cannot certify safety under every possible condition.

1. **Significance:** Validation adds dimensions that test-set accuracy ignores. A model achieving 95 percent accuracy on a static test set may miss a 100 ms latency target on its slowest requests, underperform a required subgroup threshold such as a maximum five percentage-point performance gap across demographic groups, or lose more than ten percentage points of accuracy under common image-quality problems such as blur and glare. Each unchecked dimension is a deployment risk that compounds silently in production.
2. **Distinction:** Evaluation measures model behavior using chosen datasets, metrics, and experiments, which may include out-of-distribution and subgroup tests. Validation assembles the broader evidence needed for a specified deployment decision, including system latency, throughput, robustness, operational readiness, and cost.
3. **Common pitfall:** A frequent misconception is that validation is “one more test.” In reality, it is a multi-dimensional gate: a model can pass the accuracy test and still fail validation because it violates latency, subgroup reliability, or cost constraints that accuracy alone does not measure.

Evaluation and validation address different questions. Evaluation characterizes model behavior with selected datasets, metrics, and experiments. Validation asks whether the accumulated model and system evidence supports a specified use in a specified environment, including edge cases, distribution shifts, operational constraints, and unusual input conditions. Together these processes establish the evidence base required for deployment decisions and define validation as a risk-management discipline.

3.6.1 Evaluation metrics and thresholds

Effective evaluation begins with metrics that align with problem definition objectives. For our DR screening system, standard classification metrics like accuracy prove insufficient. The illustrative requirements set sensitivity above 90 percent to limit missed referable disease and specificity above 80 percent to avoid overwhelming referral systems with false positives. Actual thresholds depend on the intended use, population, and regulatory evidence.



Production validation is a gate: field sensitivity below the required floor fails deployment.

Beyond aggregate metrics, stratified evaluation reveals performance variations across patient subgroups. A model achieving 94 percent overall accuracy might drop below 80 percent for patients with specific comorbidities, particular age groups, or images captured under certain lighting conditions. These disparities, invisible in aggregate metrics, become critical in production where every patient deserves reliable predictions. Chapter 12 provides systematic treatment of these evaluation methodologies.

Evaluation must also address **calibration**:²⁴ whether a predicted 80 percent confidence corresponds to 80 percent observed correctness. Poorly calibrated models undermine clinical trust even when accuracy metrics appear strong. Clinicians relying on confidence scores for triage decisions need those scores to reflect true uncertainty.

3.6.2 Offline and online evaluation

The validation gate begins offline and then moves progressively closer to production. *Offline evaluation* on held-out test sets establishes baseline performance but cannot guarantee production behavior. *Online evaluation* deploys models in controlled production conditions through **staged rollout**:²⁵ shadow mode runs the model without serving its predictions, canary deployment exposes a limited share of traffic, and A/B testing uses randomized allocation to compare outcomes against a baseline. The required traffic for either method depends on the effect and evidence sought. This staged rollout meaning is separate from the cross-tier compression pattern called progressive deployment in chapter 2; here the emphasis is validation risk, not producing smaller model variants.

Each validation stage adds different, overlapping evidence, as table 3.4 summarizes.

Table 3.4: Staged Validation Coverage: The stages provide complementary, overlapping evidence rather than partitioning failures into exclusive classes.

Validation stage	Typical evidence it adds
Offline evaluation	Algorithmic behavior
Shadow mode	Integration behavior
Canary deployment	Limited-traffic behavior
A/B testing	Comparative outcomes

Teams should plan for this staged validation workflow from the beginning, because retrofitting staged rollout to an already-deployed system proves more difficult than building it into the original deployment architecture.

3.6.3 Production-condition validation

After staged evaluation establishes basic behavior, production-condition validation tests whether the model survives the specific environment named during problem definition. This process reveals failure modes that standard evaluation cannot detect, and its three checks move outward in scope: from the sites a model meets at deployment, to the individual inputs it sees at inference, to the slow drift of the world it operates in.

External and multisite validation address the site-level question of whether the model has learned generalizable patterns or overfit to characteristics specific to its development sources. A DR model trained primarily on images from high-quality research cameras must demonstrate robust performance on images from the diverse equipment deployed across clinic networks. Validation datasets should include images from equipment manufacturers, lighting conditions, and operator skill levels representative of actual deployment contexts.

Even within a single validated site, individual inputs vary, so robustness testing narrows the lens to the single inference, subjecting models to realistic perturbations and edge cases. For image-based systems, this includes testing with varying brightness, contrast, focus quality, and partial occlusions. In our DR example, teams discover that models optimized for research-quality images may fail on images captured by technicians with minimal training, requiring preprocessing pipelines that normalize image quality before inference.

24 | Calibration: A calibrated model’s predicted probabilities match observed frequencies: 80 percent confidence should correspond to 80 percent correctness. Calibration is distinct from accuracy, and the systems consequence for the DR system is severe: clinicians use confidence scores for triage decisions, so a miscalibrated model that assigns 90 percent confidence to uncertain cases can misdirect clinical workflows more dangerously than a less accurate but well-calibrated alternative. Platt scaling and temperature scaling can improve posttraining calibration, but their effectiveness must be evaluated on representative data.

25 | Staged rollout: Shadow mode runs the new model in parallel, logging predictions without serving them. Canary deployment (named after coal-mine canaries) then exposes a limited, configurable share of traffic and increases it gradually if metrics hold. Coverage depends on traffic, observability, duration, and the failure modes exercised, so staged rollout reduces risk without guaranteeing that it catches a fixed share of production issues. Chapter 14 details implementation strategies.

A model that passes both checks today can still fail next year, which is why temporal validation extends the horizon, assessing whether models maintain performance over time. Data distributions shift as patient populations change, equipment ages, and clinical practices evolve. Models validated only on historical data may degrade unexpectedly when deployed. This includes data or covariate drift when the input distribution changes, and **concept drift** when the relationship between inputs and labels changes; both motivate the continuous monitoring discussed in section 3.8.

3.6.4 Regulatory validation

Healthcare AI systems face device- and pathway-specific validation requirements. When FDA premarket review is required, the safety and performance evidence must be appropriate to the device and its regulatory pathway; clinical data are required for some submissions, not all.²⁶

Domain-specific validation goes beyond regulatory compliance to address stakeholder requirements. Clinical validation studies in our DR example involve deploying the system alongside expert graders and comparing predictions against ground truth established by consensus panels of ophthalmologists. These studies must demonstrate comparable accuracy and acceptable failure modes: systems that *fail safely* (referring uncertain cases to specialists) receive more clinical trust than those that *fail silently*.

Human factors validation assesses how clinicians interact with system predictions and whether the overall workflow achieves intended outcomes. A technically accurate model that clinicians distrust or misuse fails to deliver clinical value. Validation studies should measure end-to-end workflow outcomes (clinician confidence, referral appropriateness, and patient satisfaction) alongside model performance metrics.

3.6.5 Deployment readiness

Successful validation produces the evidence package for deployment decisions: documentation covering performance across relevant metrics and subgroups, characterization of failure modes and their frequencies, validated preprocessing and inference pipelines, and evidence of regulatory compliance where required. The transition from validation to deployment represents a decision point where teams assess whether accumulated evidence supports production release. This decision balances technical performance metrics, operational readiness, regulatory status, and organizational capacity for monitoring and maintenance. Incomplete validation creates deployment risks that compound throughout the system lifecycle.

Validation failures drive model architecture revisions, training data augmentation, and preprocessing pipeline improvements. Validation successes establish the performance baselines and monitoring thresholds that guide production operations. Once the model evidence clears this gate, deployment still depends on system readiness: validated preprocessing and serving paths, observability, rollback procedures, operational ownership, and any required regulatory authorization. The central issue shifts from laboratory performance to whether the complete system can operate reliably in its target environment.

3.7 Deployment and Integration

A model that passes every validation test in the lab still faces its hardest exam when it meets the real world. Consider the DR scenario: a validated model must now run on tablets in rural clinics with intermittent connectivity, integrate with hospital information systems it was never tested against, and produce results that clinicians trust enough to act on. The scenario's latency and connectivity assumptions rule out cloud round-trips. Deployment is where the abstract constraints specified during problem definition become concrete engineering requirements. In iron-law terms, the serving path determines data movement, execution efficiency, and fixed overhead together; which term binds varies by archetype. ResNet-50-class vision workloads may use batching to improve throughput, DLRM-class recommendation workloads emphasize interactive latency and embedding access, and TinyML-class audio workloads operate under severe energy budgets (table 3.2). Chapter 14 covers the operational aspects of deployment and maintenance in depth.

26 | **FDA AI/ML regulation:** FDA marketing authorization requires evidence appropriate to the device, intended use, risk, and regulatory pathway. The FDA's 2021 AI/ML SaMD Action Plan identified lifecycle management, real-world performance monitoring, and pre-determined change control planning as central issues for adaptive medical-device software (U.S. Food and Drug Administration 2021). These concerns make model versions, training provenance, monitoring, and change documentation important workflow artifacts without imposing one identical audit artifact at every lifecycle stage.

3.7.1 Deployment requirements

The requirements for deployment stem from both the technical specifications of the model and the operational constraints of its intended environment. In our DR example, the model must operate in rural clinics with limited computational resources and intermittent internet connectivity, while automated quality checks flag poor-quality images for recapture. It must fit into the existing clinical workflow, requiring rapid, interpretable results that assist healthcare providers without causing disruption.

Napkin Math 3.3: Cloud vs. edge deployment economics

Problem: A production model processes about 760,000 billable screening images per month across 500 clinics, assuming daily operation and one processed image per patient after local selection and quality checks. Should the deployment use a cloud inference endpoint or edge inference on an on-premises server?

Option A: Cloud inference.

- Model runs on centralized GPU servers
- Inference cost: ~\$0.01/image (cloud GPU time + API overhead)
- Annual inference cost across 500 clinics, 50 patients/day, 1 billable image per patient, 365 days/year, and \$0.01/image is \$91,250/year
- Plus: An assumed connectivity and network-operations allocation for uploading 5 MB per billable image = ~\$45,000/year
- **Total:** ~\$136,250/year operational cost
- Risk: 200 ms+ exceeds the scenario latency target; connectivity outages halt screening

Option B: Edge deployment on a clinic device.

- One-time hardware: one \$500/device per clinic across 500 clinics requires \$250,000 capital expense
- Inference cost: ~\$0.001/image (assumed all-in variable operating allocation)
- Annual cost: approximately \$25,000/year maintenance plus approximately \$9,125/year variable inference cost, for about \$34,125/year
- **Total:** \$250,000 upfront + ~\$34,125/year
- Benefit: latency below 50 ms; works offline; much lower per-inference cost


Math:

- Annual savings: \$136,250/year (cloud) – \$34,125/year (edge) = \$102,125/year
- Payback: \$250,000 (hardware) ÷ \$102,125/year = ~2.4 years

Systems insight: Under these scenario assumptions, edge deployment pays back in ~2.4 years and allows inference during connectivity outages, yet it requires tighter model optimization (must fit in edge memory) and more complex update pipelines. The deployment paradigm selected during Problem Definition determines whether the edge option is even viable.

The comparison makes the deployment trade-off concrete. In this scenario, bandwidth, latency, connectivity, and workflow constraints favor edge deployment and therefore set tight model size, latency, and memory budgets that the systematic compression techniques in chapter 10 must satisfy. Once compressed, the model must be served efficiently under latency and throughput constraints; chapter 13 addresses the serving infrastructure that bridges optimized models and production traffic.

Integration with existing systems poses additional challenges. The ML system must interface with hospital information systems (HIS) for accessing patient records and storing results. HIS integration for a probabilistic ML model differs fundamentally from integrating a deterministic sensor: rather than storing a single crisp value like a blood pressure reading, the system must communicate



confidence scores and uncertainty estimates that the HIS can display and that clinicians can act on. The interface contract must encode a clinically validated routing policy, including which cases are referred automatically and which require physician review; no confidence threshold is inherently safe outside that policy and its evidence. Privacy regulations add a further ML-specific constraint: HIPAA and analogous frameworks govern not only secure storage and transmission but also whether production inferences can be retained, re-associated with patient records, and legally fed back into a model-update workflow. If the system cannot retain certain outputs, the improvement path must rely on other approved evidence, making privacy compliance a constraint on model maintenance rather than only on data transit. Chapter 14 details operational considerations that apply to these deployments.

3.7.2 Pilot to full deployment

Deployment proceeds through phases that progressively expose the system to real-world complexity, because each phase adds different evidence. Simulated environments can catch integration issues before any real users are affected. Pilot sites reveal variability that simulation may miss, including equipment differences, operator skill levels, and patient populations. Broader deployment exposes a longer tail of image artifacts, lighting conditions, and rare clinical presentations. Conditions that were absent or sparse in development data warrant explicit monitoring and uncertainty-aware handling, but their rarity alone does not prove that every prediction is unreliable.

Scaling across multiple sites adds to these challenges. Each clinic presents unique constraints (different imaging equipment, varying network reliability, diverse operator expertise levels, and distinct workflow patterns), creating data quality inconsistencies that can require preprocessing adjustments not exposed during the pilot. The deployment paradigm itself constrains solutions: edge deployment can reduce network latency but imposes strict model complexity limits, while cloud deployment enables flexibility but introduces network latency that may violate clinical workflow requirements.

Successful deployment requires more than technical optimization. Clinician trust depends on model calibration and usable explanations, not just aggregate accuracy: a clinician who cannot interpret the system's uncertainty cannot know when additional review is warranted. **Human-in-the-loop routing** applies a validated policy to send designated cases to a specialist rather than presenting every prediction as actionable. Automated image-quality checks, calibrated routing rules, and stress testing for peak volumes all contribute evidence for reliable operation; none substitutes for clinical workflow validation.

Managing improvements across distributed deployments requires centralized version control and automated update pipelines. Deployment feedback (usability concerns, performance regressions, integration surprises) shapes the monitoring strategies that keep the system healthy over time. Deployment is not an endpoint but a transition into continuous operations, where the system's behavior must be watched as carefully as any patient it screens.

3.8 Monitoring and Maintenance

Consider an illustrative case six months after a DR screening system launches. A clinic upgrades its fundus cameras, and the new equipment produces images with different color profiles. The model's sensitivity then drops at that site because the pixel distributions it learned during training no longer match the images it receives. No code changed. The data drifted beyond the training envelope, and the model degraded silently. *ML systems can degrade through data drift even when their artifacts remain untouched.* This possibility means that deployment is not the end of the lifecycle but the beginning of an ongoing operational phase. Monitoring provides the statistical telemetry to detect degradation; maintenance ensures the system evolves in response. Chapter 14 develops these operational practices in full.

In this illustrative scenario, monitoring tracks performance across clinics to detect whether changing patient demographics, camera technology, or equipment degradation affects accuracy. Adding a new imaging modality such as optical coherence tomography would be a product change requiring new data, validation, workflow integration, and any applicable regulatory review rather than routine maintenance. Three feedback pathways guide update decisions: performance evidence can motivate targeted data collection, data-quality findings can prompt preparation changes, and confirmed

model degradation can justify retraining or another intervention. Drift thresholds initiate review; they do not select the remedy automatically.

3.8.1 Production monitoring

Monitoring must serve two audiences simultaneously: technical teams tracking system health metrics and clinical staff needing actionable insights. Initial deployment may reveal blind spots invisible during laboratory validation.²⁷ Clinics with older equipment may show accuracy decreases. Specific patient subgroups, such as those with proliferative retinopathy or cataracts complicating the fundus image, may trigger higher error rates. These discoveries drive targeted data collection and architectural improvements.

A DR screening system, where missed referable disease can delay care and contribute to avoidable vision loss, demands continuous operational monitoring plus periodic performance evaluation when labels arrive. Teams establish quantitative thresholds for latency, accuracy, and data distribution stability. Lightweight statistical tests such as population stability index (PSI) and Kolmogorov-Smirnov (KS) tests²⁸ can trigger responses ranging from on-call alerts to retraining review; chapter 14 develops the monitoring pipelines around these tests.

A production DR system can track four metric categories at different timescales. The following thresholds illustrate one scenario policy; clinical and operational teams must validate them for the intended product, population, and deployment environment.

- **Model performance metrics** (requiring ground truth, available with delay): sensitivity (target above 90 percent, alert if seven-day rolling average drops below 88 percent), specificity (target above 80 percent, alert if it drops below 78 percent), and subgroup performance (alert if any demographic drops more than 5 percentage points below baseline).
- **Proxy metrics** (available immediately, without ground truth): prediction confidence distribution (alert if mean confidence drops more than 10 percent relative to baseline), referral rate (alert if rate changes more than 15 percent from baseline), and image quality rejection rate (alert if more than 20 percent of images fail quality checks).
- **Operational metrics**: Inference latency (p95 below 50 ms, alert if above 100 ms), throughput (alert if queue depth exceeds 50 images), and error rate (alert if more than 0.1 percent of requests fail).
- **Data stability metrics**: Feature and prediction distributions compared with a baseline, with alerts when recent traffic moves outside the expected range.

The hierarchy matters: operational metrics can surface immediate problems, proxy metrics can flag possible model issues without waiting for ground truth, and performance metrics often arrive later because they require labeled data.

3.8.2 Maintenance at scale

Model updates require careful validation and controlled rollouts. Teams employ A/B testing frameworks to evaluate updates and implement rollback mechanisms²⁹ that address issues quickly. ML systems must account for data evolution that can affect behavior independently of code changes.

In the illustrative scenario, scaling from pilot sites to hundreds of clinics increases monitoring complexity. The resulting log volume depends on request rates, sampling, payload sizes, and retention policies as well as the number of clinics. The monitoring infrastructure must track both global metrics and site-specific behaviors, maintain **data lineage**,³⁰ the metadata trail linking production logs to data, code, and model versions, where required for regulatory compliance, and correlate production issues with training experiments for root cause analysis.


Proactive maintenance closes the lifecycle loop: operational signals identify potential problems, review determines whether newly validated data or another change is needed, and production evidence feeds back to refine problem definitions, data-quality standards, and architectural decisions. When retraining is selected, scheduled or evidence-triggered pipelines can incorporate approved data and return the candidate through validation. The patterns underlying these dynamics (why constraints propagate, why feedback operates at multiple timescales, and why system-level behavior diverges from component-level behavior) are the subject of section 3.9.

²⁷ | **Lab-to-clinic performance gap:** Medical AI systems can experience substantial performance changes when camera models, image quality, patient populations, or operator workflows differ from development data; the gap arises because training data cannot capture the full diversity of production conditions. FDA has emphasized total product lifecycle oversight and real-world performance monitoring for AI/ML-enabled medical devices, and device submissions may need evidence appropriate to the product's intended use and risk. For ML systems engineers, this means monitoring infrastructure should be a deployment prerequisite, not a postlaunch addition.

²⁸ | **Population stability index (PSI) and Kolmogorov-Smirnov (KS) test:** Two lightweight statistical methods for detecting distribution drift; PSI bins features and computes divergence; thresholds such as 0.1 and 0.2 are operating conventions, not universal significance levels. The KS test measures maximum distance between empirical cumulative distributions. Both are cheap enough for frequent monitoring, but a detected input shift does not by itself prove an accuracy change; labels or other outcome evidence are needed. Chapter 14 covers drift detection pipelines in depth.

²⁹ | **ML rollback complexity:** An ML model's validity is coupled to the data distribution on which it was trained, not just its code. "Data evolution" means a simple rollback restores a model artifact but cannot restore the past data environment, creating a temporal state mismatch. Even a rapid rollback is therefore a mitigation tactic rather than a true system restore, as the stale model's performance on live data is not guaranteed.

30 | Data lineage: The automated recording of meta-data linking each clinic's production logs to the exact data, code, and model version that generated them. Without this explicit trail, correlating a site-specific accuracy drop with a training experiment can require manual forensic analysis that delays root-cause identification.



An illustrative doubling scenario visualizes late-discovery sensitivity.

3.9 Systems Thinking

The DR scenario showed the lifecycle acting as a coupled system: bandwidth, latency, connectivity, and workflow constraints favored edge deployment; edge deployment constrained model size; and model size reshaped preprocessing. Three structural patterns explain that cascade. Recognizing them transforms reactive debugging about deployment failure into proactive design that surfaces downstream constraints early.

3.9.1 Constraint propagation principle

The DR scenario illustrated constraint propagation repeatedly: bandwidth, latency, connectivity, and workflow constraints favored edge deployment, which constrained model size and reshaped data preprocessing. Each decision narrowed the feasible design space for dependent stages. This narrowing gives the pattern its name.

 **Definition 3.3:** The constraint propagation principle

The constraint propagation principle states that a constraint discovered late in the ML lifecycle can force rework in affected earlier stages. Actual correction cost depends on which artifacts and decisions must change. This chapter's $2^{N_{\text{stage}}-1}$ rule is an illustrative sensitivity scenario, not an empirical cost law.

1. **Significance:** A 100 ms latency target discovered at deployment (stage 5) may propagate backward to constrain model size (stage 3: algorithm complexity O), dataset requirements (stage 2: dataset size D), and problem definition (stage 1: what accuracy is achievable). Rework depends on which earlier decisions relied on the missing constraint. Within the iron law, a deployment constraint on L_{lat} or R_{peak} can redefine the feasible region for O , D_{vol} , and η_{hw} .
2. **Distinction:** Unlike modular decomposition (which encourages independent optimization of each component), this principle mandates end-to-end reasoning: optimizing accuracy in isolation may produce a model that is infeasible to deploy, making the “local maximum” in accuracy a “global minimum” in system viability.
3. **Common pitfall:** A frequent misconception is that deployment is “the last step.” In reality, the deployment environment is the day-one constraint: its latency budget, memory capacity, and power envelope define the boundaries of every upstream decision.

Propagation operates bidirectionally, creating dynamic constraint networks rather than linear dependencies. When rural clinic deployment reveals tight bandwidth limitations, teams may redesign the pipeline to transmit compact outputs after local inference. If the system instead compresses model inputs, the model architecture and training strategy must account for the resulting data degradation. Understanding these cascading relationships enables teams to make architectural decisions that accommodate rather than fight against systemic constraints.

The constraint propagation principle formalizes what experienced ML engineers know intuitively: decisions made in ignorance of downstream constraints can create technical debt.³¹ The stage interface specification (table 3.3) operationalizes this principle by making constraints explicit at each stage boundary, aligning with the model, data, and infrastructure contract practices discussed in chapter 14. Those contracts enable earlier detection of constraints. When propagation occurs specifically through data quality failures, the resulting pattern is known as a **data cascade**: a chain of downstream failures triggered by bad data (Sambasivan et al. 2021). Chapter 4 formalizes this failure mode and traces how it unfolds stage by stage.

3.9.2 Multi-scale feedback

ML systems succeed by orchestrating feedback loops across multiple timescales, each serving a different purpose. One illustrative operating plan for the DR scenario uses minute-level checks to catch a misconfigured camera before it produces a shift's worth of unusable images; daily reviews for proxy shifts such as referral-rate, confidence, rejection-rate, or site-specific camera changes; weekly

31 | ML technical debt: Sculley et al. (2015) identify ML-specific debt mechanisms such as *entanglement* (changing one feature affects all others because the model learned joint distributions), *hidden feedback loops* (predictions influence future training data), and *undeclared consumers* (downstream systems depending on outputs without contracts). Since ML code is often only a small part of a production system, the surrounding configuration, pipelines, and infrastructure can allow this debt to accumulate silently.

aggregation of labeled accuracy statistics and drift tests when ground truth is available; monthly analysis of population coverage; and quarterly review of whether the architecture still meets clinical needs. Actual cadences depend on label delay, product risk, traffic, and the cost of acting on each signal.

The temporal structure of these feedback loops reflects the inherent dynamics of ML systems. Rapid loops enable quick correction of operational issues—a clinic’s misconfigured camera can be detected and corrected within minutes. Slower loops enable strategic adaptation; recognizing that population demographic shifts require expanded training data takes months of monitoring to detect reliably. This multi-scale approach prevents both reactionary changes (over-responding to daily fluctuations) and sluggish adaptation (under-responding to meaningful trends). Concretely, fast iteration is not just a productivity metric; it is a systems feature that expands opportunities to discover better architectures and hyperparameters.

3.9.3 Emergent complexity and resource trade-offs

Complex systems produce **emergent behaviors** invisible when analyzing individual components. In a multisite DR deployment, individual clinics could show stable aggregate performance while system-wide analysis detects degradation affecting specific demographic groups—patterns invisible in single-site monitoring but critical for equitable healthcare delivery. ML systems can experience probabilistic degradation through data drift and bias amplification, while both ML and conventional distributed systems can also fail through deterministic cascades such as server crashes or resource exhaustion. Probabilistic degradation may lack the obvious error signals that trigger traditional incident response.



Checkpoint 3.3: The cost of late discovery

Apply the constraint propagation principle to this scenario:


A team discovers during monitoring (Stage 6) that their DR model fails for patients over 70 years old. This demographic requirement should have been specified at Problem Definition (Stage 1).

- ☐ Cost multiplier: Calculate the relative cost multiplier using the formula from the constraint propagation principle definition.
- ☐ Affected stages: Identify which intermediate stages may need to be revisited to fix this issue.
- ☐ Prevention rule: State the design rule that would have prevented this late-stage rework.

Resource optimization introduces multi-dimensional trade-offs that also arise in conventional software, while ML adds learned statistical behavior to them. An accuracy improvement might require increasing the model size, forcing deployment onto more powerful hardware; when multiplied across many clinics, that incremental accuracy gain translates into capital expenditure. These trade-offs manifest the power wall and memory wall from chapter 2: edge deployment can reduce network latency but constrains model complexity; cloud deployment enables flexibility but introduces network latency that may violate workflow requirements. When we trace these relationships across the system, we can make strategic architectural decisions rather than optimize components in isolation.

Together, these three patterns (constraint propagation, multi-scale feedback, and emergent complexity with its attendant resource trade-offs) define the engineering discipline that transforms ML development from ad hoc experimentation into systematic practice. A late-discovery scenario tests the most consequential pattern.

In the chapter’s illustrative model, a discovery at monitoring, the final stage, carries the largest assigned multiplier. In practice, the fix propagates only through affected stages. These principles predict specific failure modes; the fallacies and pitfalls in section 3.10 capture the most common ways teams violate them.



An illustrative monitoring plan spans minutes to quarters across five orders of magnitude.

3.10 Fallacies and Pitfalls

ML workflows introduce counterintuitive complexities that lead teams to apply familiar software patterns to structurally different problems. These fallacies and pitfalls capture errors that waste development cycles, cause production failures, and create technical debt that compounds as systems scale.

Fallacy: *ML development can follow traditional software workflows without modification.*

Engineers assume waterfall or standard agile processes will work for ML projects without modification. ML adds learned behavior, statistical evaluation, and data feedback to established software-engineering concerns (table 3.1). Workflows that treat requirements as fixed and all testing as binary pass/fail cannot accommodate iterative experimentation in which problem definitions evolve through exploration. Practitioner surveys often identify data work as a major claim on practitioners' time (section 3.1.1), and rigid phase gates can prevent teams from revisiting data, model, and deployment assumptions.

Pitfall: *Treating data preparation as a one-time preprocessing step.*

Teams assume they can “finish” data preparation and move on to modeling. In production, data distributions may shift over time. The two-pipeline architecture in figure 3.1 shows data and model pipelines running in parallel with continuous feedback, not sequentially. As section 3.4 establishes, data quality decisions cascade through model training, validation, and deployment. Data quality issues are a common source of production ML failures. Recommendation systems, fraud models, and clinical models may all need feature, label, or preprocessing updates as the world changes, and unchecked drift can degrade accuracy or cause larger failures under training-serving skew. Without continuous validation, drift may be discovered only after users or operators notice degraded behavior; teams that build data validation pipelines from the start can detect drift earlier and trigger update review.

Fallacy: *Passing model evaluation means the system is ready for deployment.*

Engineers treat the model development pipeline as the entire workflow, assuming strong evaluation metrics mean the system is complete. This is single-axis thinking. The two-pipeline architecture in figure 3.1 exposes the blind spot: data-pipeline feedback loops, deployment integration, and production monitoring remain unaddressed. The diabetic retinopathy screening case study (section 3.2.1) demonstrates the gap: model evaluation on curated data did not establish behavior across clinic equipment, operator workflows, and patient populations. Evaluation can characterize many dimensions of model behavior, but production readiness requires additional system evidence, including data freshness, preprocessing consistency, latency under load, observability, and failure recovery. The constraint propagation principle (section 3.9.1) explains why late discoveries can require broader correction effort. Teams that equate favorable model metrics with deployment readiness underestimate the integration work.

Pitfall: *Scaling data collection before checking marginal model value.*

Teams assume that scaling dataset size is the most reliable path to accuracy gains, treating data collection as a monotonically beneficial investment. In practice, additional data may provide little benefit once the target distribution is sufficiently covered, while labeling, storage, and preprocessing costs continue to grow. The feedback loops in figure 3.1 illustrate why: model performance depends on the interaction between data quality, model capacity, and deployment conditions, not data volume alone. A smaller dataset with careful label quality control and balanced class representation can outperform a larger dataset with noisy labels and skewed distributions. The Data Collection and Preparation stage (section 3.4) establishes that data quality decisions cascade through every subsequent stage. Teams should compare the marginal value of cleaning existing data against collecting more data.

Fallacy: *Skipping validation stages accelerates delivery.*

Teams assume cutting validation time ships faster. In production, the multi-stage validation process exists because each stage catches different failure modes (section 3.6). Skipping shadow mode testing can expose integration issues such as latency spikes only after launch. Bypassing canary deployment

can turn localized model failures into broad user-facing incidents. Postdeployment fixes can be more expensive than catching issues during validation because they combine incident response, rollback, root-cause analysis, data repair, and renewed validation. A team that “saves” time by skipping validation may spend substantially more time on emergency remediation. Organizations investing in systematic validation infrastructure can catch production-condition failures earlier.

Pitfall: *Deferring deployment paradigm selection until after model development.*

Teams assume they can “figure out deployment later” and focus first on model accuracy. In production, deployment paradigm (Cloud, Edge, Mobile, TinyML) is not a late-stage detail; it is a binding constraint shaping every preceding stage (table 3.3). Suppose a team develops a 2 GB ensemble model before learning that its TinyML target has 256 KB of memory. The resulting cascade requires revisiting Data Collection, Model Development, and Evaluation. The chapter’s illustrative constraint-propagation model assigns a stage-5 discovery $2^4 = 16\times$ the stage-1 cost. Teams that defer paradigm selection create avoidable iteration cycles and schedule risk. The paradigm determines what can be built, not merely where it runs.

3.11 Summary

The lifecycle is a feedback loop, not a checklist. The data pipeline transforms raw inputs through collection, ingestion, analysis, labeling, validation, and preparation into ML-ready datasets. The model development pipeline takes these datasets through training, evaluation, validation, and deployment to create production systems. With the full chapter as context, the feedback arrows in figure 3.1 carry the chapter’s central meaning: each one represents a lesson learned in production flowing back to strengthen earlier stages, making data and model feedback explicit in the development cycle.

Understanding this framework explains why machine learning systems require specialized additions to established software-engineering practices. ML workflows add probabilistic optimization, learned statistical behavior, and data-dependent feedback loops. The iron law supplies one quantitative lens: decisions across the lifecycle jointly change data movement, operation count, hardware efficiency, and fixed latency in $(T = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}})$, while production evidence feeds constraint violations back into review. This perspective recognizes that success emerges not from perfecting individual stages in isolation, but from understanding how data quality affects model performance, how deployment constraints shape training strategies, and how production insights inform each subsequent development iteration.

Three patterns define the engineering discipline of the ML lifecycle. Many practitioners report data collection, cleaning, labeling, validation, or preparation as major time sinks, while model development is only one part of the lifecycle. Production-ready systems also require repeated iteration cycles across data, model, and infrastructure stages, so investment in data engineering can have high leverage when data quality is a major source of rework. Finally, late constraints can increase rework: the chapter’s illustrative constraint-propagation model assigns a constraint discovered at stage N_{stage} roughly $2^{N_{\text{stage}}-1}$ times the stage-1 correction cost, and it assigns a deployment paradigm mismatch discovered at stage 5 a $16\times$ cost multiplier. Actual rework depends on which stages are affected. Early constraint discovery does not require freezing the workflow. It establishes an operating envelope through latency, memory, power, safety, and validation obligations while data and model choices continue to iterate within those bounds.



Key Takeaways: See the whole map first

- **The lifecycle is a loop, not a checklist:** Data and model pipelines advance in parallel, but production feedback is what makes them a system. Monitoring, validation, and retraining send lessons from deployment back into collection, labeling, architecture, and infrastructure decisions.
- **Late constraints can increase rework:** The illustrative model assigns a deployment limit found at stage N_{stage} roughly $2^{N_{\text{stage}}-1}$ times the stage-1 correction cost. The stage-5 mismatch’s $16\times$ multiplier shows why requirements should flow backward early.

- **Iteration velocity expands search opportunity:** Under the worked assumptions, a light-weight model starting 5 percentage points behind reaches the shared 99 percent ceiling because shorter cycles permit more assumed improvements. Workflow speed alone cannot guarantee model quality.
- **Interfaces make feedback actionable:** Stage contracts define inputs, outputs, and quality invariants so that data, model, and deployment teams can detect violations before integration. Without explicit contracts, each stage can optimize locally while the system fails globally.
- **Production speaks on different clocks:** Real-time inference monitoring, batch update reviews, and quarterly architectural reviews answer different failure modes. Treating all feedback as one loop either reacts too slowly to drift or churns expensive workflows without signal.
- **Workflow carries constraints through time:** Problem definition records targets, data collection shapes evidence and byte volume, model development changes operations and reuse, deployment realizes the complete serving path, and monitoring sends violations back into review. The lifecycle is data-algorithm-machine coupling unfolding over time.


This workflow framework turns ad hoc ML experimentation into a more disciplined engineering practice. By understanding how data pipelines and model development interact through feedback loops, teams can identify integration risks earlier and allocate resources more deliberately. The constraint propagation principle shows why systematic workflow management is risk mitigation rather than bureaucratic overhead.

Seen as a checklist, the lifecycle is a sequence of stages to clear in order. Seen correctly, it is a loop. A constraint discovered at deployment does not stay there; it may force rework in dependent model and data decisions. A workflow is therefore a coupled system rather than a pipeline, the same coupling the D·A·M taxonomy describes in *space*, now unfolding in *time*. Optimizing one stage in isolation moves the failure instead of removing it, which is why the discipline is to see the whole map before touching any single part of it.

What's Next: From blueprint to fuel

The workflow framework established here provides the blueprint for every technical chapter that follows. An engine without fuel, however, is only a heavy block of metal; for an ML system, that fuel is data. Chapter 4 therefore develops the first D·A·M axis, Data. Part II then builds the core system (neural network computation, model architectures, frameworks, training), Part III optimizes it for deployment (data selection, model compression, hardware acceleration, benchmarking), and Part IV deploys and operates it in production (model serving, operational practices, responsible engineering). Each chapter assumes familiarity with where its techniques fit within this complete workflow, building on the systematic perspective developed here.

Data Engineering




- 4.1 Dataset Compilation
- 4.2 Physics of Data
- 4.3 Four Pillars Framework
- 4.4 Data Acquisition
- 4.5 Data Pipeline Architecture
- 4.6 Systematic Data Processing
- 4.7 Data Labeling
- 4.8 Storage Architecture
- 4.9 Operational Data Health
- 4.10 Fallacies and Pitfalls
- 4.11 Summary

Purpose

Why does training data function like source code in a machine learning system?

In conventional software, programmers write logic that computers execute. In machine learning, programmers write optimization procedures that extract operational logic from data. This inversion makes training data resemble source code: changing it can change the learned system when the model is retrained, even if no traditional code changes. Subtle labeling inconsistencies can induce behavioral inconsistencies; missing edge cases can leave corresponding failures unmeasured; historical biases can propagate into learned behavior. A model trained only on that dataset cannot infer task-relevant evidence absent from it, and systematic label errors can shape what it learns. Unlike traditional source code, which sits inert until a programmer modifies it, the operating distribution can shift as the world changes, potentially changing production performance even when nothing in the codebase has been touched. Data engineering can consume substantial project effort because the work is consequential. Every decision in the data pipeline (what to collect, how to label, when to filter, how to split) propagates forward to constrain model architecture, training dynamics, and deployment viability. Data engineering is therefore not preprocessing but programming in a different language, one where quality control, versioning, and monitoring determine whether the compiled system works today and continues working tomorrow. In D-A-M terms, the pipeline itself is a data-machine co-design problem: however well curated the data, the algorithm can only learn as fast as the machine can deliver it.





Learning Objectives

- Explain data as source code and trace how data cascades propagate through ML systems
- Calculate data gravity, feeding-tax, and storage-bandwidth costs for moving or serving datasets
- Evaluate acquisition strategies against coverage, quality, labeling cost, governance, and deployment constraints
- Design ingestion and validation pipelines for batch, streaming, ETL, and ELT workloads
- Implement idempotent transformations, lineage, and drift checks to preserve training-serving consistency
- Select labeling, storage, file format, versioning, and feature-store designs for ML lifecycle needs
- Diagnose data debt and production pipeline failures using quality, reliability, scalability, and governance evidence

4.1 Dataset Compilation

THE workflow becomes concrete in the data pipeline: raw inputs pass through collection, ingestion, analysis, labeling, validation, and preparation before they become ML-ready datasets. The effort breakdown reported in the 2016 Crowdflower survey (CrowdFlower 2016) explains why that pipeline needs its own systems treatment: 60 percent of respondents selected cleaning and organizing data as their most time-consuming task, and 19 percent selected collecting datasets. The data axis of the D·A·M taxonomy becomes real only as infrastructure: acquisition systems, validation checks, labeling workflows, storage layouts, and governance.



Definition 4.1: Data engineering

Data engineering is the infrastructure layer that manages the lifecycle of data from source to model, encompassing acquisition, transformation, storage, and governance.

1. **Significance:** Its critical function is ensuring training-serving consistency, preventing silent degradation by decoupling the model from the volatility of raw data. Within the iron law, it governs bytes moved (D_{vol}), while dataset composition and quality determine whether the dataset (D) remains representative of the target distribution.
2. **Distinction:** Unlike data science, which focuses on inference and insight, data engineering addresses the scalability and reliability of the data pipeline.
3. **Common pitfall:** A frequent misconception is that data engineering is “data cleaning.” In reality, it is dataset compilation: transforming raw, noisy observations into an optimized binary that the model consumes.

The **dataset compilation** analogy makes the infrastructure concrete. Just as a software compiler transforms source code through a series of increasingly refined representations (tokens, abstract syntax trees, intermediate representations, machine code), a data pipeline transforms raw, noisy observations into training-ready numeric tensors through analogous stages. Filtering corrupted records, outliers, and irrelevant features corresponds to *dead code elimination*: stripping material that contributes nothing to the learned representation. Augmentation, which synthetically expands limited examples by rotating images, pitch-shifting audio, or injecting noise, mirrors *loop unrolling*, exposing the model to more variations of the underlying pattern without collecting new data. **Deduplication** plays the role of *common subexpression elimination*, identifying and merging duplicate records that would otherwise bias gradient estimates and waste compute. Schema validation, enforcing strict types and ranges on every record, is the data pipeline’s *type checker*, rejecting malformed inputs before they crash the “runtime” of model training.

The engineering implication is that datasets must be versioned (like git), unit-tested (data quality checks), and debugged. *Deleting a row of training data can change the learned artifact, just as deleting a line of code can change a compiled binary; retraining is the corresponding recompilation step.* Compilation also forces a trust boundary between dataset partitions. The training set is allowed to shape parameters; the validation set is allowed to shape modeling and pipeline choices; the test set is reserved for estimating generalization after those choices have been made. Leakage occurs when information crosses those boundaries: duplicate examples appearing in multiple splits, augmented variants of the same source record landing on both sides of the split, user records from the same household appearing in both training and test, or time-derived features computed using future observations. For the keyword-spotting case study used throughout this chapter, this means speaker-independent and time-aware splits are not bookkeeping details; they are the difference between measuring memorization of familiar voices and measuring performance on the deployment population.

This compilation metaphor establishes the engineering mindset that runs through the chapter. A compiler has distinct phases (lexing, parsing, optimization, code generation), and our dataset compiler has phases too: acquisition, ingestion, processing, labeling, storage, and ongoing maintenance. A **four pillars framework** of Quality, Reliability, Scalability, and Governance organizes design decisions across all phases. A keyword spotting (KWS) case study illustrates each stage under extreme resource constraints, where every byte and operation matters.

Sound pipeline design begins with the physical properties that constrain each stage. Just as a civil engineer must understand soil mechanics before designing foundations, a data engineer must understand the physics of data movement and information density before making pipeline decisions. These physical constraints impose hard boundaries that no amount of clever software can circumvent.

4.2 Physics of Data

The “data as code” metaphor captures *what* data does (determines system behavior) but not *why* moving it is so expensive. The physics of data explains why data systems must treat data as a physical substance with measurable properties. Just as diverse materials have density and viscosity, datasets have task-relevant signal and data gravity.

4.2.1 Data gravity


Data gravity is the cost of movement. It is a function of volume (D_{vol}) and network bandwidth (BW). The time to move a petabyte dataset across a 10 Gbps link is fixed by physics ($T = D_{\text{vol}}/\text{BW} \approx 9.3$ days); even a 100 Gbps dedicated link leaves transfer time and egress cost large enough to shape the architecture. This gravity dictates architecture: because moving 1 PB to the compute is slow and expensive, the compute often must move to the data. This explains the rise of “Data Lakehouse” architectures¹ (Zaharia et al. 2021) where processing engines such as Spark and Presto operate over shared object storage. In contrast, **data mesh** (Dehghani 2022) proposes decentralizing ownership to manage this scale organizationally, treating data as a product owned by domain teams.

4.2.2 Task-relevant signal

Task-relevant signal is a heuristic for the useful information a dataset contributes to a particular task, not Shannon entropy computed from raw bytes. A dataset of 1 million identical images has high gravity but little additional task-relevant information beyond the repeated image. A dataset of 10,000 diverse edge cases may have lower gravity but greater task-relevant signal. Let this heuristic measure useful signal per byte and data gravity capture movement cost. Their ratio represents a dataset’s return on movement cost:

$$\text{Data Selection Gain} \propto \frac{\text{Task-Relevant Signal}}{\text{Data Gravity}} \quad (4.1)$$

The heuristic favors datasets whose marginal task signal justifies their movement cost, so deduplication and active learning improve the selection decision when they remove redundant bytes or prioritize informative examples.



Selection gain peaks where signal density (entropy) is high and movement cost (gravity) low.

¹ | **Data lakehouse:** Combines data lake storage (cheap, schema-less) with warehouse query semantics (ACID transactions, schema enforcement) using transactional table layers such as Delta Lake. For ML workloads, the lakehouse reduces the extract, transform, load (ETL) copy between lake and warehouse, enabling direct feature computation on the storage layer where data already resides – a direct response to data gravity, since repeated petabyte-scale copies increase the D_{vol}/BW cost (Armbrust et al. 2020; Zaharia et al. 2021).

4.2.3 The feeding problem: Flow rate and the “feeding tax”

Data gravity establishes the cost of moving the entire mass; **the feeding problem** establishes the cost of delivering it. We analyze this as a **flow rate** problem: the struggle to saturate a high-throughput machine from a low-bandwidth data source.

According to the iron law, the system is only as fast as its slowest term. If a high-throughput accelerator running an image model can process 1,843 img/s, but the storage pipeline delivers only 250 MB/s, the expensive silicon sits idle. We quantify this as **the feeding tax**: the wall-clock time lost to I/O wait, which directly reduces the system efficiency (η_{hw}) term. For a standard cloud volume, the feeding tax can exceed 77.5 percent, meaning the accelerator spends the majority of its time waiting for bits. This tax transforms the data pipeline from a simple storage concern into the primary regulator of the system’s duty cycle. Feeding the reference accelerator in this example often requires 1.1 GB/s transfer rates, forcing the shift from traditional file systems to the specialized storage architectures developed in section 4.8.1. These physical properties also carry an energy cost: as data moves farther from the processor, movement increasingly dominates the budget.

🔍 Systems Perspective 4.1: The energy-movement invariant

The iron law in section 1.7 and the memory-wall analysis in section 2.4 imply a data-engineering principle: *moving one 32-bit value costs roughly 173× to 10,810.8× as much energy as the table’s FP32 multiply*, with the multiplier rising as the data leaves the chip. Later optimization chapters exploit the same cost gradient inside model and hardware design; here, the question is how much energy the information flow from the external world already consumes before the model computes. Table 4.1 makes that cost gradient explicit across the memory hierarchy.

Table 4.1: Energy Cost of Moving vs. Computing on a 32-Bit Value: Approximate energy per event at each stage of the data-movement hierarchy, with multipliers normalized to one 32-bit FP32 multiply. All rows use the same 32-bit width, so the ratios compare like-sized events. They do not by themselves prove which term dominates a complete workload; total energy also depends on the number of arithmetic operations and transfers.

Operation	Energy (pJ)	Relative Cost
32-bit FP Multiply	3.7 pJ/FLOP	1×
DRAM Memory Access (32-bit)	640 pJ	173×
Local SSD Access (32-bit)	4,000 pJ	1,081.1×
Network Transfer (32-bit)	40,000 pJ	10,810.8×

The cost gradient quantified in table 4.1 explains why locality can be a high-leverage systems optimization. Each avoided transfer removes its associated movement energy, although the total benefit depends on access counts and reuse.

Data has physical mass. Pruning 50 percent of training data through deduplication does more than save disk space; it reduces work in data-intensive stages of the training lifecycle. Data selection is therefore a high-leverage tool when data movement and processing dominate the workload.

The energy argument and the signal-density heuristic point to the same lever from two directions: effective data engineering maximizes the Data Selection Gain defined in equation 4.1. “Data Cleaning” is more than basic hygiene: it is **signal-to-noise engineering**. Deduplication can remove redundant mass while preserving task-relevant signal, directly increasing the ratio. Active learning can prioritize informative examples over redundant ones, thereby increasing useful information per byte. We optimize this ratio to ensure our storage and compute budgets are spent on signal, not noise.

These principles operate at the level of individual files and batches. At data center scale, the cost of moving data compounds into a constraint that makes large datasets so expensive to transfer that compute must relocate to the data rather than the reverse. Transferring a petabyte dataset between data centers puts a dollar figure on the constraint.

Napkin Math 4.1: The physics of data gravity

Problem: A 1 PB training dataset resides in a US East data center, while a remote accelerator pod is available in US West. Is it faster to move the data or to move the compute?

Math:

1. **Network bandwidth:** $100 \text{ Gb/s} \div 8 = 12.5 \text{ GB/s}$.
2. **Transfer time:** $1,000,000 \text{ GB} \div 12.5 \text{ GB/s} = 80,000 \text{ seconds}$, approximately 22 h.
3. **Cost:** $1,000,000 \text{ GB} \times \$0.02/\text{GB} = \$20,000$. (Baseline: AWS data transfer out pricing, 2024.)

Systems insight: If training takes less than 22 h, data transfer takes longer than training. If training costs less than \$20,000 (approximately 5,000 TPUv4-hours), bandwidth costs more than compute. For petabyte-scale data, code moves to data; for gigabyte-scale data, data moves to code.

These physical constraints govern every decision in production data pipelines. Before moving on, check the fundamental intuitions that will recur throughout the pipeline.

Checkpoint 4.1: The physics of data

Data engineering is governed by physical costs. Check your intuition:

- ☐ **Data gravity:** why do petabyte-scale datasets force compute to move to the data?
- ☐ **energy-movement invariant:** why does moving a byte of data cost orders of magnitude more energy than processing it?
- ☐ **Task-relevant signal:** why can a smaller, diverse dataset be more valuable than a massive, redundant one?

These physical properties impose hard constraints on every pipeline decision: where to store data, how to transform it, and when to move computation rather than bytes. Physics alone, however, does not prevent failures; it merely defines the boundaries within which engineering decisions must be made. A team that understands data gravity perfectly can still build a brittle pipeline if quality checks are ad hoc, error handling is absent, or governance is an afterthought. Translating physical constraints into reliable practice requires a systematic framework that organizes design decisions across every pipeline stage.

4.3 Four Pillars Framework

A credit scoring model rejects every applicant from a region because an upstream team changed a ZIP code field from integer to string. A medical imaging model degrades silently for months because camera hardware changed at a partner hospital. A fraud detection system misses a new attack vector because its training data was six months stale. Each failure traces to a different root cause (schema drift, distribution shift, data staleness), yet all share a common pattern: ad hoc data engineering decisions that interacted in ways no one anticipated until deployment. These cascading failures motivate a four-pillar framework organized around quality, reliability, scalability, and governance.

4.3.1 Data cascades

Machine learning systems face a distinctive failure pattern called data cascades, where poor data quality in early stages amplifies throughout the entire pipeline (Sambasivan et al. 2021). Some invalid inputs trigger immediate software errors, but data can also pass schemas and tests while degrading learned behavior silently² until quality issues become severe enough to require expensive investigation, rework, or retraining.

Data errors rarely stay confined to their point of origin; they compound across pipeline stages. Follow the sequence of connected stages in figure 4.1 from left to right to trace how an initial data defect propagates forward into feature extraction, model optimization, and deployment decisions.

² | **Data cascades:** The failure is “silent” because it degrades model *inputs*, not model *code*—corrupted data can pass unit tests and appear healthy in ordinary system monitoring. Sambasivan et al. (2021) describe cascades as often invisible and delayed: flawed data practices may surface only after downstream evaluation, deployment, or user-facing failures reveal that the model learned from the wrong signal. Remediation then requires tracing the issue back through data collection, labeling, feature engineering, and evaluation decisions, with some teams restarting or abandoning affected work.




Figure 4.1: Data Quality Cascades: Upstream defects along the seven-stage workflow timeline cast failure arcs directly into late-stage evaluation and deployment. Early flaws in problem framing and collection compound through downstream processing, forcing costly project restarts (dashed red return paths) when defects finally surface in production. Source: (Sambasivan et al. 2021).



Definition 4.2: Data cascade

Data cascade is an ML systems failure mode in which upstream data quality problems propagate through collection, labeling, feature engineering, training, evaluation, and deployment, amplifying into downstream model failures or user harm.

1. **Significance:** The remediation cost scales with the number of downstream artifacts that consumed the corrupted data: feature tables must be regenerated, models retrained, evaluation metrics recomputed, and deployed systems rolled back or patched. A single schema or labeling defect can therefore invalidate an entire training run and all experiments derived from it, turning a local data error into a full-pipeline rework event.
2. **Distinction:** Unlike an isolated data quality bug, a data cascade is defined by propagation and amplification. The initial defect may be small, but each pipeline stage treats its input as trustworthy and converts the defect into new derived state, making the root cause harder to observe as the system moves farther from collection.
3. **Common pitfall:** A frequent misconception is that data cascades are caught by ordinary software tests. In reality, corrupted data can satisfy schemas, pass unit tests, and produce successful training jobs while teaching the model the wrong signal; preventing cascades requires lineage, validation, contracts, and monitoring tied to model behavior.

This feedback loop creates an insidious operational failure mode: silent degradation. Unlike conventional software crashes that trigger stack traces and alert pages, data-induced model errors produce syntactically valid outputs with degraded semantic accuracy. By the time downstream metrics reflect the failure, corrupted predictions have already contaminated historical data stores, forcing expensive rollbacks and data re-ingestion.



Example 4.1: The pipeline jungle

Scenario: A production credit scoring model suddenly begins rejecting all applicants from a specific geographic region.

Diagnosis: An upstream database team altered the `zip_code` field schema from `integer` to `string` ("02139") to support international codes without updating downstream data contracts. A legacy downstream transform still coerced the string back to an integer before serializing it, stripping the leading zero ("2139") and causing the model to treat a familiar postal code as an unobserved, high-risk category.

Systems lesson: Unenforced feature boundaries create pipeline jungles where upstream schema modifications cause silent downstream model failures. Versioned data contracts with strict schema and range validation prevent cross-system data corruption.

4.3.2 Four foundational pillars

To prevent silent failures from compounding across the pipeline, data infrastructure must balance four competing operational priorities. Observe how figure 4.2 organizes these priorities into foundational columns spanning input integrity, deployment stability, execution performance, and data governance.




Figure 4.2: The Four Pillars of Data Engineering: Quality, Reliability, Scalability, and Governance form the foundational framework for ML data systems. Each pillar contributes essential capabilities (solid arrows), while trade-offs between pillars (dashed lines) require careful balancing: validation overhead affects throughput, consistency constraints limit distributed scale, privacy requirements impact performance, and bias mitigation may reduce available training data.

These four operational foundations establish the trade-off space for every architectural decision in this chapter. Prioritizing scalability without quality validation or reliability contracts can increase throughput while allowing bad or inconsistent data to propagate. Conversely, exhaustive governance controls can add storage and computation overhead that constrains training iteration velocity.

Reliability asks whether the same enrollment pipeline keeps working when the device, network, or user behavior is imperfect. A pipeline that produces excellent wake-word examples in a laboratory but fails during intermittent connectivity, battery pressure, or microphone glitches delivers no value in the user's home. Error handling, retries, local buffering, and graceful degradation turn the quality rule into a usable system: the device can request another utterance, defer upload until connectivity returns, or preserve a known-good enrollment rather than silently accepting corrupted audio.

Scalability asks whether that same decision survives growth. A manual review policy that works for a thousand recordings collapses when the product expands to millions of users, dozens of languages, and long-tail acoustic conditions. The system must scale validation, storage, labeling, and retraining without letting infrastructure cost grow faster than the value of better coverage. The limiting resource therefore depends on the workload.

Lighthouse 4.1 shows where the scalability pillar bites hardest: the wall is memory capacity, and partitioning embedding tables across machines becomes a first-class design concern rather than an afterthought. Governance then defines the boundaries within which quality, reliability, and scalability may operate. For the KWS example, governance determines whether raw voice recordings may leave the device, how long enrollment audio can be retained, which consent record authorizes use, and what documentation proves that the dataset covers relevant accents without exposing private speech. A perfectly scalable, reliable, high-quality pipeline that violates the GDPR or perpetuates demographic biases creates liability rather than value. Dataset documentation practices such as data statements make part of that governance visible by recording provenance, intended use, collection conditions, and coverage needed for bias analysis and scientific comparison (Bender and Friedman 2018).

When ML systems exhibit failures, the four pillars provide a diagnostic lens for identifying root causes. Gradual accuracy degradation points to quality: data drift has shifted the serving distribution

away from training, or label quality has degraded as annotator pools change. Intermittent pipeline failures point to reliability: error handling, retry logic, or resource controls are missing under peak load. Training that takes too long despite adequate hardware points to scalability: a single-threaded transformation, unpartitioned shuffle, or slow storage tier prevents parallel resources from being used. Compliance gaps discovered during audits point to governance debt: lineage tracking is incomplete, access controls are stale, or retention policies have not kept pace with regulatory changes.

 Lighthouse 4.1: DLRM recommendation lighthouse

Recommendation systems built in the DLRM style exemplify the scalability challenge of modern data engineering. Modern recommendation workloads bifurcate pipeline resource demands: continuous signals stream through dense matrix pipelines, while categorical IDs demand terabyte-scale distributed lookup infrastructure. Table 4.2 maps these operational stress points across data ingestion, embedding table aggregation, and feature interaction stages.

Table 4.2: DLRM Scalability Profile: Recommendation models stress data engineering along memory capacity and sparse-access bandwidth rather than compute.

Property	Value	System Implication
Data scale	High-cardinality user/item IDs	Embedding and lookup tables can outgrow one machine’s memory.
Constraint	Memory Capacity	The tables no longer fit on one machine and must be partitioned.
Bottleneck	Sparse Access	Random lookups stress memory bandwidth more than compute.

At ingestion, high-cardinality categorical IDs remain compact records in a throughput-bound stream. The network- and memory-capacity bottleneck emerges later, when training or serving must fetch sparse embeddings from terabytes of partitioned table state. Mitigating this asymmetry requires the feature-store caching and table-partitioning strategies developed later in the pipeline lifecycle.

The most insidious failures span multiple pillars. Features that differ between training and serving implicate both quality (the values are wrong) and reliability (the computation is inconsistent). A privacy-motivated deletion policy can also create quality gaps if the retained data no longer covers the deployment population. Diagnosing such cross-pillar failures requires checking consistency contracts, comparing feature distributions across environments, and tracing transformation lineage, all techniques examined in detail throughout this chapter. Production diagnosis should therefore check data infrastructure alongside model behavior.

4.3.3 KWS case study

KWS systems provide an ideal case study for applying our four-pillar framework to real-world data engineering challenges. These systems power voice-activated devices like smartphones and smart speakers, detecting specific wake words such as “OK, Google” or “Alexa” within continuous audio streams while operating under strict resource constraints.³

The broad operational goals of data engineering translate into concrete techniques, tools, and validation checks. For the voice assistant exchange illustrated in figure 4.3, these mechanisms include deduplication, schema contracts, prefetching, and lineage tracking, mapped to their corresponding architectural pillars.

The four pillars translate directly into engineering constraints for the KWS system. The core problem is deceptively simple: detect specific keywords amidst ambient sounds and other spoken words, with high accuracy, low latency, and minimal false activations, on devices with severely limited computational resources. A well-specified problem definition identifies the desired keywords, the envisioned application, and the deployment scenario. The objectives that follow must balance competing requirements: performance targets of 98 percent accuracy in keyword detection with latency under 200 ms, alongside resource constraints demanding minimal power consumption and model sizes optimized for available device memory.

³ | **Voice Match enrollment:** Repeating “OK Google” creates a micro-scale data pipeline: quality filters noisy samples, reliability completes enrollment, scalability fits the model in always-on system-on-chip (SoC) memory, and governance controls storage, processing, and retention. Data engineering applies wherever data determines system behavior.






Figure 4.3: Keyword Spotting in Action: An Amazon Echo Dot exchange shows the user saying “Alexa,” the device acknowledging the wake word, and the user issuing “dim the lights.” A lightweight, always-on wake-word detector listens continuously and triggers the main voice assistant only after the keyword is spoken.

Napkin Math 4.2: False positive targets

Problem: An always-on KWS system must tolerate at most one false wake-up per month. With one-second classification windows running around the clock, how strict does the per-window false positive rate need to be?

Variables:

- **Duty cycle:** Always-on (24 hours/day).
- **Window size:** One-second classification windows.
- **Windows per month:** One window per second, 24 hours/day, over 30 days gives 2,592,000 windows/month.

Math:

- **False positive rate (FPR):** 1 tolerated false wake divided by the monthly window count gives approximately 3.9×10^{-7}
- **Nonkeyword rejection requirement:** 99.99996 percent rejection of nonkeyword windows.

Systems insight: Aggregate accuracy (for example, “99 percent accuracy”) is insufficient here. Evaluation must report false accepts per hour (FA/Hr) together with the corresponding false-rejection behavior.

Success metrics for KWS extend beyond simple accuracy to include true positive rate (correctly identified keywords relative to all spoken keywords), false positive rate (nonkeywords incorrectly identified as keywords), and detection/error trade-off curves that compare false accepts per hour against false rejection rate on streaming audio representative of real-world deployment, as demonstrated by Nayak et al. (2022). Of these metrics, the false positive rate deserves particular attention for always-on systems. Because KWS listens continuously, every second of every day, even a seemingly negligible false positive rate compounds across millions of evaluation windows. A quick calculation shows how strict that requirement becomes.

Operational metrics further track response time (keyword utterance to system response) and power consumption (average power used during keyword detection), and stakeholder priorities create additional tension around those metrics. Device manufacturers prioritize low power consumption, software developers emphasize ease of integration, and end users demand accuracy and responsiveness. Balancing these competing requirements shapes system architecture decisions throughout development.

Embedded device constraints impose hard boundaries on these architectural choices. Memory limitations require extremely lightweight models, often in the tens-of-kilobytes range, to fit in the always-on island of the SoC;⁴ this constraint covers only model weights, and preprocessing code must also fit within tight memory bounds. Limited computational capabilities (often a few hundred MHz of clock speed) demand aggressive model optimization. Most embedded devices run on batteries, so KWS systems target sub-milliwatt power consumption during continuous listening.

⁴ | **System-on-chip (SoC) always-on island:** Modern system-on-chip designs partition power domains so a low-power “always-on” island (typically achieving sub-milliwatt draw) monitors for wake triggers while the main processor sleeps. The critical constraint is that this island must hold both the model weights *and* the audio preprocessing code within its dedicated SRAM—a split budget that forces KWS architectures to optimize for total footprint, not just parameter count.

Devices must also function across diverse deployment scenarios ranging from quiet bedrooms to noisy industrial settings.

Data quality and diversity ultimately determine whether these constraints can be met. The dataset must capture demographic diversity (speakers with various accents, ages, and genders) to ensure broad recognition. Keyword variations require attention since people pronounce wake words differently, and background noise diversity proves essential for training models that perform across real-world scenarios from quiet environments to noisy conditions. Once a prototype system is developed, iterative feedback and refinement keep the system aligned with objectives as deployment scenarios evolve, requiring testing in real-world conditions and systematic refinement based on observed failure patterns.

4.3.4 KWS design space

KWS accuracy, false-wake tolerance, latency budget, energy budget, and memory limits create a multi-dimensional design space where data engineering choices cascade through system performance. Table 4.3 quantifies key trade-offs, enabling principled decisions rather than ad-hoc selection. One row uses mel-frequency cepstral coefficients (MFCCs): compact speech-frequency features whose coefficient count controls feature size, compute cost, and acoustic detail; section 4.6 shows how they are extracted.

Table 4.3: KWS Data Engineering Design Space: Illustrative scenario inputs that price each design choice along four axes, namely quality, latency, cost, and memory; the entries are not measurements from one system and should not be combined as additive effects. Higher sampling rates double raw-audio storage and processing, while more training examples multiply labeling and storage requirements. Actual quality and latency must be measured on the target pipeline.

Design Choice	Quality Impact	Latency Impact	Cost Impact	Memory Impact
16 kHz vs. 8 kHz sampling	+2–4% accuracy	2× input-processing workload	2× raw-audio storage	2× feature size
13 vs. 40 MFCC coefficients	+3–5% accuracy	3× feature compute	Minimal	3× feature memory
1M vs. 10M training examples	+5–8% accuracy	10× training time	10× labeling cost	10× storage
Clean vs. noisy training data	+10–15% real-world	Minimal	3× collection cost	Minimal
Local vs. cloud inference	up to 2% accuracy risk	10 ms vs. 100 ms	\$0/query vs. \$0.001/query	64 KB vs. cloud-scale
Synthetic vs. real augmentation	+3–5% robustness	Minimal	10× cheaper	Minimal

Always-on keyword spotting operates under milliwatt power envelopes and strict tens-of-kilobytes SRAM budgets. Choosing an architecture requires balancing feature resolution against memory footprint; table 4.3 catalogs these trade-offs across model size, inference latency, and detection accuracy for representative edge deployments.



Napkin Math 4.3: Optimizing the KWS design space

Scenario: A KWS system for a smart speaker has these constraints:

- **Target:** 98 percent accuracy, fewer than 1 false wake/month
- **Budget:** \$150K total data engineering budget
- **Memory:** 64 KB model size limit (always-on island)
- **Timeline:** 6 months to production

Step 1: Apply constraints to eliminate options.

For this scenario, the team chooses two conservative options:

- 13 MFCC coefficients to reduce feature-state memory and compute
- Local inference to avoid network latency and preserve offline operation

The 64 KB model-size limit does not by itself prove that either alternative is infeasible; the deployed model, preprocessing state, runtime, and network stack require separate footprint measurements.

Step 2: Calculate budget allocation.

The \$150K budget splits across three cost categories at the unit rates established in section 4.4.2:

- **Labeling** (~60 percent): \$90K available
- **Storage/Processing** (~25 percent): \$37.5K
- **Governance/Other** (~15 percent): \$22.5K

At \$0.10/label with 20 percent review overhead: $\$90K \div \$0.12/\text{label} = 750K$ labeled examples

This yields roughly 0.75M labeled examples, just below the 1M anchor in our design space. The 1M-to-10M row in table 4.3 should therefore be read as a scaling reference rather than a direct interpolation: the real-label budget alone does not buy the full data-volume gain.

Step 3: Maximize remaining accuracy.

The quality plan has three components:

- An illustrative base-model accuracy of ~90 percent
- Partial data-volume gain from 750K real examples
- Need: higher sampling rate plus synthetic/noisy augmentation to reach the 98 percent target

Three options from the design space remain within budget and memory constraints:

- 16 kHz sampling: +2–4 percent accuracy, $2\times$ storage cost ✓ (fits budget)
- Noisy training data: +10–15 percent real-world accuracy, $3\times$ collection cost
- Synthetic augmentation: +3–5 percent robustness, $10\times$ cheaper than real data ✓

Step 4: Final configuration.

The scenario combines these choices with 13 MFCCs and the computed 750K real-label budget in table 4.4. The quality effects can overlap, so the table records a candidate configuration rather than a solved optimum.

Result: The budget calculation supports approximately 750K human-labeled examples. Reaching the 98 percent quality target and fitting the 64 KB model limit remain validation requirements for the trained and deployed system.

Systems insight: Systematic design-space analysis turns “we need more data” into a testable allocation: the stated labeling assumptions buy about 750K real examples, while synthetic volume, accuracy, and deployed footprint still require measurement.

Table 4.4 records the resulting configuration, pairing each design choice with the constraint that forces it: sampling and augmentation are bought where they fit the budget, while precision and memory are pinned by the always-on island.

Table 4.4: KWS Candidate Configuration: A scenario that uses the computed real-label budget and records the remaining sampling, augmentation, and deployment choices for measurement against the quality and memory targets.

Choice	Selection	Rationale
Sampling rate	16 kHz	+3% accuracy worth $2\times$ storage within budget
MFCC coefficients	13	Reduces feature-state memory and compute; final fit requires measurement
Training examples	750K real + illustrative synthetic augmentation	Human-label count follows from the budget; synthetic volume is a scenario choice
Data diversity	Noisy + clean mix	Critical for real-world deployment
Inference	Local, reduced precision	Reduced precision can lower footprint; the final 64 KB fit requires the compression analysis in chapter 10
Augmentation	Heavy synthetic	$10\times$ cost efficiency

With a candidate configuration selected from our design space, implementation requires combining multiple data collection approaches: preexisting corpora for the baseline, web scraping and crowdsourcing for coverage gaps, and synthetic generation for scale. The combination can improve coverage across diverse real-world conditions. Section 4.4 develops each of these acquisition strategies, their economics, and their KWS instantiations in full.

The four pillars provide the evaluative lens; the first concrete engineering decision is data provenance. Acquisition strategy determines the raw material that every subsequent stage refines.



Checkpoint 4.2: Four pillars framework

The four pillars provide a systems lens for every pipeline choice.

Pillars

- ☐ **Quality:** can you distinguish *correctness* from *coverage* (edge cases, subgroup coverage) and explain why both matter?
- ☐ **Reliability:** can you name concrete failure modes (schema drift, missing values, skew) that cause silent degradation rather than explicit crashes?
- ☐ **Scalability:** can you explain which parts of a pipeline become bottlenecks as data volume grows (storage bandwidth, compute, coordination), and why?
- ☐ **Governance:** can you articulate what must be tracked to make a dataset auditable (provenance, consent constraints, access controls, documentation)?

Trade-offs

- ☐ **Pillar tension:** given one pipeline change (for example, stricter validation, stronger privacy constraints, or more data sources), can you predict which pillar improves and which pillar(s) you stress?

4.4 Data Acquisition

Data acquisition begins when the team names the coverage gap the model must close. The full ImageNet database⁵ grew to about 14.2 million labeled images across 21,841 synsets, while the ImageNet Large Scale Visual Recognition Challenge used a 1,000-class subset (Deng et al. 2009; Russakovsky et al. 2015). GPT-3's training corpus used tens of terabytes of raw Common Crawl text filtered to hundreds of gigabytes, combined with curated web, books, and Wikipedia data (Brown et al. 2020). Our KWS system needs 23.4 million audio samples spanning 50 languages, but the challenge extends beyond raw volume. The model must recognize wake words across accents, microphones, rooms, ages, and background noises that no single collection method can economically cover. Acquisition strategy is therefore a sequence of gap-closing decisions: reuse what already matches deployment, collect what is missing, scrape or synthesize where scale is the binding constraint, and reject sources whose provenance or consent constraints make them unusable.

The KWS case also shows why acquisition cannot optimize one pillar at a time. Achieving 98 percent accuracy across diverse acoustic environments requires representative data spanning accents, ages, and recording conditions. Maintaining consistent detection despite device variation requires recordings from different microphones and capture paths. Supporting millions of concurrent users requires volumes that manual collection cannot economically provide. Protecting user privacy in always-listening systems constrains which recordings may be retained and how they must be anonymized. A source that improves scale but weakens governance, or improves quality but excludes important speakers, does not solve the acquisition problem.

4.4.1 Data source evaluation and selection

The choice among curated datasets, expert crowdsourcing, controlled web scraping, and synthetic generation depends on which source best closes the next deployment-distribution gap at acceptable cost, quality, and governance risk. Evaluation therefore begins with the cheapest reusable source and escalates only when the remaining gap justifies new collection or synthesis.

⁵ | **ImageNet:** The 2009 paper reported 3.2 million images across 5,247 synsets; the full database later reached about 14.2 million images across 21,841 synsets, whereas the challenge used a 1,000-class subset (Deng et al. 2009, 2024; Russakovsky et al. 2015). Its value as a benchmark is inseparable from its data engineering: Fei-Fei Li's team built labeling infrastructure that later teams could reuse. The catch is sensitivity to annotation procedures and modest distribution shifts: models tuned to ImageNet's distribution can underperform on related but shifted test sets (Recht et al. 2019; Beyer et al. 2020).

Preexisting datasets from repositories such as Kaggle, UCI (Dua and Graff 2024), and ImageNet are the first test. They offer speed and comparability when the deployment distribution resembles the benchmark enough to make reuse meaningful. For KWS, a curated speech corpus can establish the baseline model and reveal which words, languages, and acoustic conditions are already covered. Its value is not guaranteed coverage; its value is that it makes the remaining gaps measurable.

That reuse decision depends on documentation quality, which directly affects reproducibility, an ongoing crisis in machine learning research (Pineau et al. 2021; Henderson et al. 2018). Good documentation captures collection methodology, variable definitions, and baseline performance, enabling validation and replication. At scale, volume and variety compound quality challenges (Gudivada et al. 2017), requiring systematic validation pipelines rather than ad-hoc inspection.

Standard validation metrics can create an illusion of model readiness when evaluated on unrepresentative data. Shared dataset usage (figure 4.4) illustrates how training multiple models on one common dataset propagates shared biases, blind spots, and systemic limits across an entire ecosystem.




Figure 4.4: Shared Dataset Bias Propagation: Five models (A through E) all train on a single central dataset repository, reached by a path labeled shared training data. Beneath the models, five labeled arrows (bias propagation, shared limits, blind spots, common weaknesses, and systemic issues) converge on one box, limited real-world alignment. Training every model on one dataset does not produce five independent defects; it produces one correlated failure across the ecosystem.

This distribution gap represents the primary risk of static offline benchmarks. Because loss functions only optimize over observed samples, models exploit spurious artifacts present in training sets that disappear in production environments, making automated drift monitoring and continuous data validation mandatory operational requirements.

4.4.2 Scalability and cost optimization

Quality-focused data acquisition approaches face inherent scaling limitations. When scale requirements dominate, needing millions or billions of examples that manual curation cannot economically provide, web scraping and synthetic generation offer paths to massive datasets. Data-acquisition scalability requires understanding the economic models underlying different acquisition strategies: cost per labeled example, throughput limitations, and how these scale with data volume. Cost-effectiveness inverts with scale: what works at thousands of examples becomes prohibitive at millions, while high-setup-cost approaches amortize favorably at large volumes.

The per-unit economics at each stage determine which strategy dominates. Labeling a single medical image, for example, can cost orders of magnitude more than storing it for a year, a ratio that reshapes budget allocation for any team operating under fixed funding. Table 4.5 and table 4.6 provide essential context for acquisition decisions.

ML engineers should use the dated reference assumptions in table 4.5 and table 4.6 as inputs to workload-specific comparisons, not as current price quotes or directly comparable totals.

Table 4.5: Illustrative Data Engineering Costs (2024): Dated reference assumptions for workload-specific cost models, not current vendor quotes. Because the rows use different units, compare them only after scaling each rate by the quantity it prices.

Operation	Cost	Notes
Crowdsourced image label	\$0.01–0.05	Simple classification
Bounding box annotation	\$0.05–0.20	Per box, simple scenes
Expert medical label	\$50–200	Per study, radiologist
S3 storage (Standard)	\$23/TB/month	Hot storage
S3 retrieval (Glacier)	\$0.02/GB	Standard: 3-5 hours
Cloud GPU training hour	\$2–4	Cloud spot pricing
Human review hour	\$15–50	Depending on expertise

Table 4.6 extends the picture with illustrative durations for labeling, training, and serving operations.

Table 4.6: Illustrative Data Engineering Time Scales: The scenario spans about 9 orders of magnitude, with human labeling as its slowest listed stage. Actual durations depend on the dataset, task, hardware, and service configuration.

Operation	Duration	Bottleneck
Label 1M images (crowdsourced)	2–4 weeks	Annotation throughput
Train ResNet-50 on ImageNet	4–6 hours	Compute (8× A100, optimized)
Feature store lookup	1–10 ms	Network + cache

The contrast in this illustrative scenario matters: weeks for human labeling, hours for GPU training, milliseconds for serving. Here, labeling is the bottleneck. A \$100K labeling budget compares with \$64–\$192 for one 8× A100 ResNet-50 run, a 520.8×–1,562.5× ratio. Within that labeling spend, the assumed effort distribution is itself skewed: 80 percent of the work goes to 20 percent of features—the long tail of edge cases, rare categories, and quality exceptions.

All cost figures reflect approximate 2024 cloud provider rates and are intended to convey relative magnitudes rather than exact pricing.⁶ For this workload and cost model, human labor dominates the cost of a single training run. Other tasks, especially those that reuse existing labels or require repeated large-scale training, can have a different dominant term; teams should measure before deciding where to optimize.

Web scraping is the first lever on that cost structure, enabling dataset construction at scales that manual curation cannot match. Major vision datasets like ImageNet (Deng et al. 2024) and OpenImages (Kuznetsova et al. 2020) were built through systematic scraping, and large language models depend on web-scale text corpora (Groeneveld et al. 2024). Targeted scraping of domain-specific sources, such as code repositories (Chen et al. 2021), further demonstrates the approach’s versatility. However, production systems that rely on continuous scraping face pipeline reliability challenges: website structure changes break extractors, rate limiting throttles collection throughput, and dynamic content introduces inconsistencies that degrade model performance. Scraped results can also include historical images for contemporary queries, requiring systematic validation and cleaning.

Figure 4.5 shows one such result from a web scrape for “traffic light”: a historical photograph rather than a modern LED signal. If images from this context were overrepresented or correlated with a label, a model could learn spurious cues involving uniformed officers or historical streets rather than the visual properties required in its deployment environment.

This example reveals why the quality pillar cannot be satisfied by scale alone: no amount of additional scraped data removes the need for validation that detects and filters anachronistic or

⁶ | **Pricing ratios:** Archive-retrieval prices and tier ratios reflect provider policy, not a physical invariant. In this illustrative scenario, expedited retrieval costs three times the standard tier; actual prices and tier availability must be verified before deployment.



Figure 4.5: Data Source Noise: A black-and-white photograph from 1914 showing early manually operated electric traffic signals, illustrating how historical images can appear in modern web scraping results for contemporary queries. Such anachronistic content requires systematic validation and filtering to prevent spurious correlations in training data. Image reproduced in (Stromberg 2015).

contextually inappropriate content. Beyond technical quality challenges, legal and ethical constraints further bound what scraping can achieve. Not all websites permit scraping, and ongoing litigation around training data usage illustrates the legal uncertainty and potential consequences (Harvard Law School 2024). Teams must document data provenance, ensure compliance with terms of service and copyright law, and apply anonymization procedures when scraping user-generated content.

Crowdsourcing shifts the acquisition bottleneck from finding enough examples to controlling the quality of many parallel judgments. Platforms like Amazon Mechanical Turk (Amazon Web Services 2024a) demonstrated this at landmark scale with ImageNet, where distributed contributors categorized millions of images into thousands of classes (Deng et al. 2024). Crowdsourcing offers two systems advantages: scalability through parallel microtask distribution, and diversity through the range of perspectives, cultural contexts, and linguistic variations that a global contributor pool introduces. This diversity can improve generalization when it broadens coverage of the deployment population. The cost is that task design, validation, and iteration become part of the acquisition system: tasks can be adjusted dynamically based on initial results, enabling refinement of collection strategies as quality gaps emerge (Sheng and Zhang 2019). For our KWS system, MTurk-style platforms⁷ enable targeted collection of wake word samples across different demographics and environments, an approach particularly valuable for underrepresented languages or specific acoustic conditions.

Moving beyond human-generated data entirely, synthetic data generation changes the scaling constraint: examples can be generated algorithmically, but their value depends on whether the generator covers the deployment conditions that real collection would miss. This approach changes the economics of data acquisition by reducing human labor while increasing the burden on validation. The pipeline in figure 4.6 shows synthetic and historical data entering the same training workflow.

Synthetic data is particularly valuable for rare event coverage and data augmentation. Simulation environments enable controlled generation of edge cases that are impractical to collect naturally (NVIDIA 2024b). For image data, augmentation methods such as AutoAugment (Cubuk et al. 2019) and RandAugment (Cubuk et al. 2020) search over transformations that improve generalization, while broader image-augmentation practice is surveyed by Shorten and Khoshgoftaar (2019). For audio, SpecAugment masks time and frequency regions to improve speech recognition robustness (Park et al. 2019). For KWS, speech synthesis (Werchniak et al. 2021) and audio augmentation fill

⁷ **Amazon Mechanical Turk (MTurk):** A crowdsourcing platform that routes small tasks to distributed workers and can scale annotation beyond expert-only workflows. Snow et al. (2008) evaluate this pattern for natural-language annotation tasks, showing that non-expert annotations can be useful when task design and quality control are handled carefully. For wake-word audio collection, the same systems trade-off appears in domain-specific form: scale is attractive, but submissions still need acoustic checks such as signal-to-noise ratio, duration, and recording validity before they can safely enter the training set.




Figure 4.6: Synthetic Data Augmentation: Combining simulated edge cases with historical system data during training. Algorithmic generation supplements real-world observations to close coverage gaps for rare events and dangerous operating regimes without incurring proportional human collection costs. Source: (AnyLogic 2024).

the coverage gaps that remain after real collection, creating wake word variations across acoustic environments, speaker characteristics, and background conditions. The KWS case makes the coverage role of these techniques concrete.



Example 4.2: Synthetic data generation

Scenario: A keyword-spotting team lacks sufficient physical recordings of rare speaker accents, room acoustics, and background noise profiles.

Diagnosis: Pure physical collection cannot cover rare acoustic edge cases within project deadlines. Generating synthetic audio variations via pitch shifting, additive noise, and room impulse response simulation expands training set diversity without expensive field collection.

Systems lesson: Synthetic data generation acts as an automated data pipeline multiplier. Using synthetic augmentation targeted at known edge cases fills distribution coverage gaps that would otherwise require prohibitive physical dataset collection.

For our KWS system, 23.4 million audio samples spanning 50 languages demand a volume that manual collection cannot economically provide. A multi-source strategy that combines curated datasets, web scraping of video platforms and speech databases, crowdsourced collection, and synthetic generation addresses this scale requirement while maintaining coverage across acoustic environments and speaker demographics.

4.4.3 Coverage and diversity requirements

Scale alone does not guarantee reliable models. Coverage gaps in even large datasets (geographic bias, demographic underrepresentation, temporal drift, missing edge cases) cause systematic failures that aggregate metrics obscure (T. Wang et al. 2019; Oakden-Rayner et al. 2020). As figure 4.4 makes clear, multiple systems training on identical datasets inherit identical blind spots; diverse sourcing strategies are the defense against correlated failure modes.

Governance constraints further shape acquisition: privacy and health-data regulations such as GDPR and HIPAA limit what data can be collected and how (European Parliament and Council of the European Union 2016; United States Congress 1996), while ethical sourcing requires fair compensation and transparent use of human contributions. Section 15.5 examines the full governance infrastructure for production ML systems.

The diversity of sources (crowdsourced audio, synthetic waveforms, web-scraped content) creates specific challenges at the boundary where external data enters our controlled pipeline. Each source arrives in a different format, at a different cadence, with different quality guarantees, and the infrastructure that receives, validates, and routes this heterogeneous data must reconcile all of them.

4.5 Data Pipeline Architecture

In our compilation metaphor, **data pipeline architecture** is the *compiler frontend*: it parses heterogeneous raw inputs into a uniform intermediate representation that downstream stages can process reliably. Audio files from crowdsourcing platforms, synthetic waveforms from generation systems, and real-world captures from deployed devices all enter the training pipeline in different formats, and the pipeline must normalize, validate, and route them into a consistent internal representation. Production KWS inference consumes continuous audio under a separate low-latency constraint; this section follows the recorded examples used to build and update the model. Figure 4.7 maps that path across data sources, ingestion, processing, labeling, storage, and ML training.




Figure 4.7: End-to-End Data Pipeline: Heterogeneous raw sources (batch and streaming) normalize into unified storage before feeding downstream training and feature processing. The architecture decouples ingestion throughput from iterative transformation and labeling loops, all bounded by an overarching data-governance layer.

Each layer plays a specific role in the data preparation workflow. Selecting appropriate technologies requires understanding how our four framework pillars manifest at each stage. Quality requirements at one stage affect scalability constraints at another, reliability needs shape governance implementations, and the pillars interact to determine overall system effectiveness.

Data pipeline design is often constrained by storage hierarchies and I/O bandwidth, alongside CPU-side decoding and transformation. Understanding these constraints enables building efficient systems for modern ML workloads. Storage hierarchy trade-offs, ranging from high-latency object storage (ideal for archival) to low-latency in-memory stores (essential for real-time serving), and bandwidth limitations (spinning disks at 100–200 MB/s vs. RAM at 50–200 GB/s) shape every pipeline decision. Section 4.8 covers detailed storage architecture considerations.

Choosing between these design patterns requires matching workload characteristics to infrastructure capabilities. Streaming workloads demand attention to message durability (the ability to replay failed processing), ordering guarantees (what sequence is preserved, under what conditions), and geographic distribution. Batch workloads hinge on data volume relative to available memory, processing complexity, and whether computation must be distributed across machines. Single-machine tools suffice for gigabyte-scale data, but terabyte-scale processing often benefits from distributed frameworks that partition work across clusters. These layer interactions, viewed through the four-pillar lens, determine overall system effectiveness.

4.5.1 Quality through validation and monitoring

Consider a self-driving-car pipeline in which 15 percent of LiDAR point-cloud labels are misaligned by 10–20 cm—enough to place pedestrian bounding boxes on empty sidewalk. Because every record remains structurally valid, schema checks would miss the defect; statistical monitoring of label-to-sensor alignment could expose it.

Quality represents the foundation of reliable ML systems, and this example illustrates why. Pipelines implement quality through systematic validation and monitoring at every stage. Data pipeline issues represent a major source of ML failures. Schema changes breaking downstream processing, distribution drift degrading model accuracy, and data corruption silently introducing errors are concrete examples of the data-dependency and monitoring debt described by Sculley et al. (2015). These failures are insidious because they rarely cause obvious system crashes; instead, they slowly degrade model performance in ways that become apparent only after affecting users. Achieving quality therefore demands proactive monitoring and validation that catches issues before they cascade into model failures.



War Story 4.1: Microsoft Tay (2016)

Context: In March 2016, Microsoft launched Tay, a Twitter chatbot designed to engage with users online and adapt its responses from public interactions (Lee 2016).

Mechanism: Adversarial users coordinated toxic prompts and exploited a public interaction surface, including behavior that repeated user-supplied text. Microsoft’s public account describes abuse of the system but does not establish unrestricted online weight updates.

Impact: Tay began tweeting abusive, racist, and misogynistic statements within 16 hours of deployment.

Fix: Microsoft suspended the chatbot 16 hours after launch, apologized publicly, and stated that a relaunch would require stronger safeguards against misuse.

Systems lesson: Public ingestion paths require adversarial-input controls. Otherwise, a feature that accepts user-generated content also becomes a security surface, and harmful inputs can propagate through the system within hours.

Production teams implement monitoring at scale through severity-based alerting systems where different failure types trigger different response protocols. The most critical alerts indicate complete system failure: the pipeline has stopped processing entirely, showing zero throughput for more than five minutes, or a primary data source has become unavailable. These situations demand immediate attention because they halt all downstream model training or serving. More subtle degradation patterns require different detection strategies. When throughput drops to 80 percent of baseline levels, error rates climb above 5 percent, or quality metrics drift more than two standard deviations from training data characteristics, the system signals degradation requiring urgent but not immediate attention. These gradual failures often prove more dangerous than complete outages because they persist undetected for hours or days, silently corrupting model inputs and degrading prediction quality.

A recommendation system processing user interaction events at 50,000 records per second makes these severity tiers concrete. Its monitoring system tracks several interdependent signals. Instantaneous throughput alerts fire if processing drops below 40,000 records per second for more than 10 minutes, accounting for normal traffic variation while catching genuine capacity or processing problems. Each feature in the data stream has its own quality profile: if a feature like `user_age` shows null values in more than 5 percent of records when the training data contained less than 1 percent nulls, something has likely broken in the upstream data source. Duplicate detection runs on sampled data, watching for the same event appearing multiple times—a pattern that might indicate retry logic gone wrong or a database query accidentally returning the same records repeatedly.

These monitoring dimensions become particularly important when considering end-to-end latency. The system must track both whether data arrives and how long it takes to flow through the entire pipeline from the moment an event occurs to when the resulting features become available for model inference. When 95th percentile latency exceeds 30 seconds in a system with a 10-second service

level agreement, the monitoring system needs to pinpoint which pipeline stage introduced the delay: ingestion, transformation, validation, or storage.

Schema and latency alerts expose structural and timing failures; detecting a continuous-feature distribution shift requires a statistical comparison with the training baseline.

Napkin Math 4.4: Detecting drift with K-S test

Scenario: Monitoring `session_duration` distribution stability between the training baseline (P_0) and current serving distribution (P_t).

Analysis: Apply the Kolmogorov-Smirnov test to compare the empirical cumulative distribution functions:

1. **Compute CDFs:** Calculate cumulative distribution functions for both datasets.
2. **Calculate statistic ($\mathcal{D}_{KS}$):** Find the maximum absolute difference between the CDFs. Let $F_{P_0}(x)$ and $F_{P_t}(x)$ denote the empirical cumulative distribution functions of the training baseline (P_0) and current serving (P_t) datasets, evaluated at value x .

$$\mathcal{D}_{KS} = \max_x |F_{P_0}(x) - F_{P_t}(x)|$$

3. **Determine significance:** For two independent samples, compare $\mathcal{D}_{KS}$ to critical value $\mathcal{D}_{crit}$ based on training-baseline sample size n_0 and serving sample size n_t . The coefficient 1.36 is the large-sample approximation for significance level $\alpha = 0.05$.

$$\mathcal{D}_{crit} \approx 1.36 \sqrt{\frac{n_0 + n_t}{n_0 n_t}}$$

Result: With sample sizes $n_0 = n_t = 1000$, the critical value is approximately 0.061:

$$\mathcal{D}_{crit} \approx 1.36 \sqrt{(1000 + 1000) / (1000 \cdot 1000)}.$$

If we observe a maximum difference of $\mathcal{D}_{KS} = 0.08$, it exceeds the critical value, so we reject the null hypothesis and flag significant drift.

Systems insight: Statistical drift tests convert a vague distribution-shift concern into an operational trigger. The test should start an investigation or retraining workflow, not silently become another dashboard number.

Quality monitoring extends beyond simple schema validation to statistical properties that capture whether serving data resembles training data. Rather than just checking that values fall within valid ranges, production systems track rolling statistics over 24-hour windows. For numerical features like `transaction_amount` or `session_duration`, the system computes means and standard deviations continuously, then applies statistical tests like the Kolmogorov-Smirnov test⁸ to compare serving distributions against training distributions.

The K-S test detects drift in continuous features; section 4.5.3 gives the full taxonomy of distribution shifts (covariate, label, concept, and label-quality drift) with population stability index and KL divergence metrics for the degradation equation.

Categorical features require different statistical approaches. Instead of comparing means and variances, monitoring systems track category frequency distributions. When new categories appear that never existed in training data, or when existing categories shift substantially in relative frequency, the system flags potential data quality issues or genuine distribution shifts; for example, the proportion of “mobile” vs. “desktop” traffic might change by more than 20 percent. This statistical vigilance catches subtle problems that simple schema validation misses entirely: age values may remain in the valid range of 18–95, while the distribution shifts from primarily 25–45 year olds to primarily 65+ year olds, indicating the data source has changed in ways that will affect model performance.

Validation at the pipeline level encompasses multiple strategies working together. Schema validation executes synchronously as data enters the pipeline, rejecting malformed records immediately

⁸ | **Kolmogorov-Smirnov (K-S) test:** A nonparametric test that measures the maximum distance between two empirical cumulative distribution functions without assuming a parametric distribution family (Berger and Zhou 2014). In ML pipelines, the K-S test is commonly used as a univariate continuous-feature drift detector, with thresholds such as $p < 0.05$ serving as investigation triggers rather than universal failure rules. Discrete and categorical variables require tests and calibrations appropriate to their distributions.

before they can propagate downstream. Modern tools like TensorFlow Data Validation (TFDV) (Breck et al. 2019) automatically infer schemas from training data, capturing expected data types, value ranges, and presence requirements.

This synchronous validation remains simple and fast, checking properties that can be evaluated on individual records in microseconds. More sophisticated validation that requires comparing serving data against training data distributions or aggregating statistics across many records must run asynchronously to avoid blocking the ingestion pipeline. Statistical validation systems typically sample 1-10 percent of serving traffic, enough to detect meaningful shifts while avoiding the computational cost of analyzing every record. These samples accumulate in rolling windows, commonly one hour, 24 hours, and seven days, with different windows revealing different patterns. Hourly windows detect sudden shifts like a data source failing over to a backup with different characteristics, while weekly windows reveal gradual drift in user populations or behavior.

The most insidious validation challenge arises from training-serving skew: the failure mode where features are computed differently in training and serving. This typically happens when training pipelines process data in batch using one set of libraries or logic, while serving systems compute features in real time using different implementations. Even seemingly minor discrepancies (a materialized view⁹ refreshed weekly versus a complete join recomputed daily) can materially reduce production accuracy and can be difficult to diagnose because the system produces no obvious errors. We formalize this as the consistency imperative in section 4.6.1 and quantify its impact with a concrete example. Detecting training-serving skew requires infrastructure that can recompute training features on serving data for comparison, sampling raw serving data and processing it through both pipelines to measure discrepancies. Chapter 14 examines operational monitoring infrastructure for this challenge at scale.

4.5.2 Data quality as code

Just as unit tests protect software systems, data expectation tests protect ML pipelines. In **data quality as code**, teams use libraries like Great Expectations or Pandera to codify quality expectations as executable assertions. Listing 4.1 shows these checks assembled into a validation run that raises on failure.

These mechanical expectations become executable assertions that can run automatically. Continuous Integration and Continuous Deployment (CI/CD) integration runs expectations in the deployment pipeline, so violations fail deployments before bad data reaches training. A pipeline structured as data ingestion followed by data validation followed by training blocks deployment when validation detects anomalies like age values of 150, triggering alerts for investigation.

Executable expectations catch properties a suite can encode; the remaining question is what valid-looking data can still get wrong.

🔍 Systems Perspective 4.2: Mechanical vs. semantic quality

In traditional software, many checks are *mechanical*. A null dereference is a runtime error in languages with null references, while integer-overflow behavior depends on the language and type.

In ML systems, data quality has a second, softer dimension: *semantic* quality.

- **Mechanical check:** “Is age an integer?” (Yes/No).
- **Semantic check:** “Is the age distribution shifting?” (Probabilistic).

A dataset can be mechanically perfect (no nulls, correct types) but semantically broken (for example, all users are suddenly 25 years old due to a default value change). Robust ML systems must validate both the container (mechanical) and the content (semantic).

Once validation becomes code, expectation suites become versioned artifacts alongside training code. When training code changes, expectation updates keep data contracts evolving with it. This coupling reduces the risk of silent divergence where code assumes data properties that the upstream pipeline no longer provides.

⁹ | **Materialized view:** A database optimization that precomputes and caches query results as a physical table. For ML systems, the risk is structural: when a materialized view refreshes on a different schedule in training and serving environments, the feature values the model trained on can diverge from what it receives at inference, degrading accuracy without producing an execution error.

Listing 4.1: Data Quality Assertions: A Great Expectations suite checks age range, identifier presence and uniqueness, and allowed country codes, then raises an error when validation fails.

```
import great_expectations as gx

# Create a data context
context = gx.get_context()

# Define an expectation suite as executable quality contract
suite = gx.ExpectationSuite(name="user_data_quality")
suite = context.suites.add(suite)

# Range validation: prevents physiologically impossible values
suite.add_expectation(
    gx.expectations.ExpectColumnValuesToBeBetween(
        column="age", min_value=0, max_value=120
    )
)

# Null detection: ensures primary key integrity for joins
suite.add_expectation(
    gx.expectations.ExpectColumnValuesToNotBeNull(column="user_id")
)

# Uniqueness: prevents duplicate training examples
suite.add_expectation(
    gx.expectations.ExpectColumnValuesToBeUnique(column="user_id")
)

# Categorical validation: detects unexpected values from upstream changes
suite.add_expectation(
    gx.expectations.ExpectColumnDistinctValuesToBeInSet(
        column="country_code",
        value_set=["US", "CA", "UK", "DE", "FR"],
    )
)

# Link the suite to a preconfigured Batch Definition and run validation
# (the Batch Definition connects GX to the training_users data asset)
validation_definition = gx.ValidationDefinition(
    name="training_users_validation",
    data=batch_definition,
    suite=suite,
)
validation_definition = context.validation_definitions.add(
    validation_definition
)
results = validation_definition.run()
if not results.success:
    raise ValueError(f"Data quality check failed: {results}")
```

These checks catch many schema-level production data issues before they reach training, including missing values, invalid ranges, type errors, and contract violations. The remaining issues require runtime monitoring and outcome checks, because semantic quality problems often emerge only in the full production data stream.

4.5.3 Data drift detection and response

ML models rest on the assumption that production data resembles training data. When this assumption breaks through statistical shifts rather than an explicit contract violation, model performance can change silently without obvious errors or system failures. The validation and monitoring techniques in section 4.5.1 can reveal both sudden and gradual problems; data drift detection focuses specifically on changes in distributions that may alter model behavior over time. Detecting, investigating, and responding to these changes requires sustained monitoring effort, making drift a core data engineer-

ing responsibility rather than an optional advanced topic. Chapter 14 builds on this foundation with operational response orchestration, outcome validation, and evidence-based retraining pipelines at scale.

Section B.2.2 formalizes the divergence metrics used to make the divergence term $\mathcal{D}(P_t \| P_0)$ of the degradation equation (equation 1.3) actionable. The **population stability index (PSI)** quantifies changes in categorical or binned feature distributions, while **Kullback-Leibler (KL) divergence** measures an information-theoretic divergence between probability distributions. The two metrics encode different distributional assumptions and scales, and neither by itself supplies a universal cutoff. Teams calibrate metric-specific thresholds against a deployment baseline and monitor them over time. Crossing a threshold signals that the live distribution has moved far enough to warrant investigation and outcome checks; input divergence alone does not establish whether accuracy improved or degraded, or by how much.

Understanding the three core types of drift enables targeted detection and response strategies. Each type manifests differently in production systems and requires distinct monitoring approaches.

The first case, **covariate shift**, changes the input distribution while preserving the relationship between features and labels: $p(x)$ changes but $p(y | x)$ stays the same. A medical imaging system trained on one camera model might later receive production images from a different manufacturer. The disease-image relationship remains unchanged, but pixel values shift because sensor characteristics, color calibration, or image processing pipelines differ. Detection therefore focuses on input feature distributions, using metrics such as PSI or KL divergence.

The second case, **label shift**, changes the output distribution while preserving the relationship between labels and features: $p(y)$ changes but $p(x | y)$ stays the same. Disease prevalence might change seasonally while symptoms remain consistent predictors of each disease. A recommendation system might see the same pattern when new product categories launch, changing the relative frequency of user preferences without altering what makes products appealing within each category. Detection can often begin without ground truth labels by tracking shifts in the model's prediction distribution.

The hardest case is **concept drift**: the relationship between features and labels changes, so $p(y | x)$ evolves over time (Gama et al. 2014). Medical treatment protocols change, user preferences shift as social trends evolve, and fraud patterns adapt as attackers respond to detection systems. Unlike covariate or label shift, concept drift requires ground truth labels for detection because the system must observe whether the feature-to-label relationship itself has changed.

4.5.3.1 Label quality drift

Label quality drift¹⁰ represents a meta-level shift distinct from the three preceding distribution shifts: the reliability of ground truth labels degrades over time even when the underlying data distributions remain stable. This drift type proves particularly insidious because standard feature distribution monitoring fails to detect it. Crowdsourced labels may degrade as annotator pools change, training materials become outdated, or labeling guidelines evolve without corresponding model updates. Automated labeling systems accumulate errors as the models powering them drift from their original operating conditions. A recommendation system using click feedback as implicit labels may see label quality degrade as user behavior becomes more exploratory, as bot traffic patterns change, or as interface modifications alter how users interact with content.

Detection requires monitoring annotation consistency rather than feature distributions. Inter-annotator agreement metrics like Cohen's kappa¹¹ (κ) provide quantitative assessment. Let p_o represent observed agreement between annotators and p_e represent agreement expected by chance. Equation 4.2 defines the statistic:

$$\kappa = \frac{p_o - p_e}{1 - p_e} \quad (4.2)$$

Monitoring κ over time windows reveals degradation trends. A medical imaging annotation project might establish a baseline $\kappa = 0.85$ (almost perfect agreement) during initial data collection, then observe decline to $\kappa = 0.72$ (substantial agreement) after six months as new annotators join without receiving equivalent domain training.

- $p(x)$
- $p(y)$
- $p(y|x)$

Each drift type traces to the distribution component that shifted.

10 | Label quality drift:

Degradation in annotation reliability over time, distinct from distribution shifts in the data itself. This drift type is invisible to standard feature monitoring because the features remain stable while the labels degrade – annotator fatigue, pool turnover, or guideline evolution silently corrupt the ground truth the model learns from. Detection requires monitoring inter-annotator agreement (κ) over rolling time windows and comparing automated labels against periodic expert audits.

11 | Cohen's kappa:

Introduced by Cohen (1960) to measure inter-rater agreement while correcting for agreement expected from the raters' class marginals, which raw percentage agreement ignores. If two independent annotators each label 90 percent of images as "not spam" in a binary task, their expected agreement is 82 percent, making raw agreement potentially misleading. The statistic is denoted κ ; interpretive bands such as the Landis-Koch categories are descriptive heuristics rather than universal data-quality thresholds (Landis and Koch 1977).

For systems with calibrated model probabilities, predictive entropy provides an additional uncertainty signal. Let p_i represent the model's probability assigned to label category i , and let $\log$ denote the natural logarithm, so entropy is measured in nats. Equation 4.3 defines this measure:

$$H_{\text{pred}} = - \sum_i p_i \log p_i \quad (4.3)$$

Rising predictive entropy indicates greater uncertainty in the model's predictive distribution, but it does not identify the cause. Harder inputs, distribution shift, calibration changes, or inconsistent supervision can produce the same signal.

Mitigation strategies depend on root cause analysis. Annotator retraining addresses systematic errors from unclear guidelines at low cost with high effectiveness. Multi-annotator voting with majority or consensus rules provides high accuracy for high-stakes domains but significantly increases annotation costs. Model-assisted labeling reduces annotator fatigue but risks introducing bias if the assisting model has its own systematic errors. Expert review sampling, where domain specialists audit a random sample of annotations, enables root cause analysis when quality decline is detected but provides medium coverage of the overall annotation stream.

Operationalizing the PSI and KL divergence metrics introduced in section B.2.2 requires connecting them to automated alerts and review workflows. Data engineering is responsible for defining domain-specific drift thresholds from baseline behavior and outcome evidence, selecting monitoring windows that can expose both sudden and gradual changes, and instrumenting pipelines to compute these metrics continuously. Chapter 14 later examines how production teams combine these signals with labeled outcomes, tiered alerts, escalation paths, cold-start monitoring, and cause-specific response orchestration.

Drift detection is one dimension of the quality pillar, focused on identifying statistical changes in data distributions over time. Detecting issues, however, is only half the challenge; the other half is ensuring systems continue operating effectively even when problems surface. The need to maintain service continuity shifts the discussion from quality monitoring to the reliability pillar.

4.5.4 Reliability through graceful degradation

Reliability ensures systems continue operating when problems occur. Pipelines face constant challenges: data sources become temporarily unavailable, network partitions separate components, upstream schema changes break parsing logic, or unexpected load spikes exhaust resources. **Graceful degradation** means handling these failures through systematic failure analysis, intelligent error handling, and automated recovery strategies that maintain service continuity even under adverse conditions.

Systematic failure mode analysis for ML data pipelines reveals predictable patterns that require specific engineering countermeasures. Data corruption failures occur when upstream systems introduce subtle format changes, encoding issues, or field value modifications that pass basic validation but corrupt model inputs. A date field switching from "YYYY-MM-DD" to "MM/DD/YYYY" format might not trigger schema validation but will break any date-based feature computation. Schema evolution¹² failures happen when source systems add fields, rename columns, or change data types without coordination, breaking downstream processing assumptions that expected specific field names or types. Resource exhaustion manifests as gradually degrading performance when data volume growth outpaces capacity planning, eventually causing pipeline failures during peak load periods.

Effective error handling strategies ensure problems are contained and recovered from systematically. Intelligent retry logic for transient errors (network interruptions or temporary service outages) requires exponential backoff strategies to avoid overwhelming recovering services. A simple linear retry that attempts reconnection every second would flood a struggling service with connection attempts, potentially preventing its recovery. Exponential backoff, retrying after one second, then two seconds, then four seconds, doubling with each attempt, gives services breathing room to recover while still maintaining persistence. Many ML systems employ **dead-letter queues (DLQs)**: separate storage for data that fails processing after multiple retry attempts. This allows for later analysis and potential reprocessing of problematic data without blocking the main pipeline (Kleppmann 2016). A pipeline processing financial transactions that encounters malformed data can route it to a dead-letter queue rather than losing critical records or halting all processing.

¹² | **Schema evolution:** This failure mode arises from a lack of contract testing between upstream data producers and downstream ML consumers. While "loud" failures like a renamed column break explicit assumptions and cause immediate pipeline crashes, "silent" failures are more dangerous. A field changing type from integer to string can pass validation but corrupt feature logic without immediate detection.

In ML systems, dead-letter queues serve dual purposes beyond failure analysis. Production teams implement systematic review of DLQ contents to identify: (1) schema violations indicating upstream changes, (2) edge case patterns the model should handle, and (3) data quality issues requiring source system fixes. For example, a fraud detection system’s DLQ revealed transactions from a new payment type the model had never seen, prompting targeted data collection and retraining rather than simply logging the failures. This transforms DLQs from passive error storage into active sources for identifying model blind spots and driving improvement.

Moving beyond ad-hoc error handling, cascade failure prevention requires circuit breaker¹³ patterns and bulkhead isolation to prevent single component failures from propagating throughout the system. When a feature computation service fails, the circuit breaker pattern stops calling that service after detecting repeated failures, preventing the caller from waiting on timeouts that would cascade into its own failure.

Automated recovery engineering extends beyond simple retry logic. Exponential backoff, jitter, load shedding, and circuit breakers reduce request pressure on struggling services while allowing transient faults to recover; simply increasing timeouts can retain resources longer and worsen overload. Multi-tier fallback systems provide degraded service when primary data sources fail: serving slightly stale cached features when real-time computation fails, or using approximate features when exact computation times out. A recommendation system unable to compute user preferences from the past 30 days might fall back to preferences from the past 90 days, providing less precise but still useful recommendations rather than failing entirely. Comprehensive alerting and escalation procedures ensure human intervention occurs when automated recovery fails, with sufficient diagnostic information captured during the failure to enable rapid debugging.

Retry logic, dead letter queues, and circuit breakers are the runtime error handlers of our dataset compiler: they catch malformed inputs without halting the entire compilation. The next question is how data enters the pipeline in the first place. The choice of ingestion pattern (batch vs. streaming; extract, transform, load (ETL) vs. extract, load, transform (ELT)) determines how quickly new data reaches the model, how much infrastructure the system requires, and how these reliability patterns are concretely deployed.

4.5.5 Data ingestion

Continuing the compilation analogy, data ingestion is the *lexer*: it reads raw source (data streams) and tokenizes them into well-formed records that the rest of the pipeline can process. A critical and often overlooked constraint in ingestion design is the **input/output (I/O) bottleneck**. Teams invest heavily in expensive GPUs, but their utilization depends entirely on whether data arrives fast enough to keep them busy. In an idealized two-stage model, step time ranges from $\max(T_{\text{compute}}, T_{\text{io}})$ with full overlap to $T_{\text{compute}} + T_{\text{io}}$ with no overlap. These bounds restate the feeding problem from section 4.2.3 through a second resource lens: there the starved term was storage bandwidth; in ingestion it is most often CPU-side decode work.

If the data pipeline cannot decode images fast enough to keep the GPU busy, the expensive accelerator sits idle. This phenomenon creates a “Choke Point” where adding more GPUs yields no speedup until the input pipeline is improved, a counterintuitive result for teams expecting linear scaling from hardware investments. This bottleneck frequently occurs in computer vision, where high-resolution JPEG decode and augmentation on the CPU can dominate the input path. For this illustrative calculation, each worker supplies 375 images/s and the accelerator consumes 3,000 images/s, so 8 workers are needed to reach the crossing. This result is a capacity match, not a hardware constant: below the crossing, the input pipeline binds; above it, additional workers do not raise throughput because the accelerator ceiling binds. Actual ResNet-50 and A100 requirements depend on image transforms, storage, batch size, host CPUs, and framework configuration.

Scaling training throughput is not simply a matter of upgrading GPU hardware; the host preprocessing pipeline must supply tensors fast enough to prevent accelerator starvation. In figure 4.8, observe the intersection between the flat GPU consumption ceiling and the sublinear CPU throughput curve as batch dimensions scale.

When the dataloader throughput falls below the GPU consumption rate, accelerator utilization drops precipitously, turning expensive compute clusters into idle memory buffers. Resolving

¹³ | **Circuit breaker:** Named for its three-state behavior – closed (normal flow), open (faults blocked), half-open (recovery probe) – after the electrical safety device that interrupts current on overload. In ML data pipelines, the circuit breaker prevents a failing feature computation service from cascading timeouts through the entire serving path: once failure count exceeds a threshold, the breaker opens and the pipeline falls back to cached or default features rather than waiting on a dead service.




Figure 4.8: The Dataloader Choke Point: Training throughput (img/s) vs. DataLoader workers. The blue CPU-supply curve rises with worker count until it reaches the illustrative input-pipeline cap; the red dashed line marks accelerator consumption capacity. The lower curve is the bottleneck: the accelerator starves to the left of the crossing, while extra workers beyond the crossing do not increase accelerator throughput.

this choke point requires multi-process worker pools, asynchronous prefetching, and hardware-accelerated decompression pipelines that bypass host CPU bottlenecks entirely.

4.5.5.1 Batch vs. streaming ingestion patterns

The choice between batch and streaming is not a preference for one architecture over another; it is a judgment about how quickly data loses value and how much infrastructure cost that freshness justifies. Batch systems buy efficiency by tolerating staleness, while streaming systems buy freshness by accepting continuous operational complexity.

Batch ingestion involves collecting data in groups or batches over a specified period before processing. This method proves appropriate when real-time data processing is not critical and data can be processed at scheduled intervals. The batch approach enables efficient use of computational resources by amortizing startup costs across large data volumes and processing when resources are available or least expensive. For example, a retail company might use batch ingestion to process daily sales data overnight, updating their ML models for inventory prediction each morning (Akidau et al. 2015). The batch job might process gigabytes of transaction data using dozens of machines for 30 minutes, then release those resources for other workloads. This scheduled processing proves far more cost-effective than maintaining always-on infrastructure, particularly when slight staleness in predictions does not affect business outcomes.

Batch processing also simplifies error handling and recovery. When a batch job fails midway, the system can retry the entire batch or resume from checkpoints without complex state management. Data scientists can inspect failed batches, understand what went wrong, and reprocess after fixes. Batch jobs can be reproducible when inputs, code, configuration, randomness, ordering-sensitive operations, and external state are controlled, which simplifies debugging and validation. These characteristics make batch ingestion attractive for ML workflows even when real-time processing is technically feasible but not required.

In contrast to this scheduled approach, **stream ingestion** processes data in real-time as it arrives, consuming events continuously rather than waiting to accumulate batches. This pattern is essential for applications requiring immediate data processing, scenarios where data loses value quickly, and systems that need to respond to events as they occur. A financial institution might use stream ingestion for real-time fraud detection, processing each transaction as it occurs to flag suspicious activity immediately before completing the transaction. The value of fraud detection drops dramatically if detection occurs hours after the fraudulent transaction completes—by then money has been transferred and accounts compromised.

However, stream processing introduces complexity that batch processing avoids. The system must handle **backpressure**, the condition where downstream systems cannot keep pace with incoming data rates. During traffic spikes, when a sudden surge produces data faster than processing capacity, the system must either buffer data (requiring memory and introducing latency), sample (losing some data), or push back to producers (potentially causing their failures). Data freshness service level agreements (SLAs) specify maximum acceptable delays between data generation and processing. Meeting a 100-millisecond SLA requires different infrastructure than meeting a one-hour SLA, affecting networking, storage, and processing architectures.

Recognizing the limitations of either approach alone, many production ML systems employ hybrid approaches that combine batch and stream ingestion. A recommendation system might use streaming ingestion for real-time user interactions to update session-based recommendations immediately, while using batch ingestion for overnight processing of user profiles and item features.



Napkin Math 4.5: The cost of real-time

Problem: A pipeline ingests 1M events/s. Which is cheaper, Batch (hourly) or Stream (sub-second) ingestion?

Physics:

1. **Throughput:** At 1M events/s and 1 KB per event, the pipeline carries 1 GB/s.
2. **Stream requirements:** To sustain 1 GB/s with less than 100 ms latency, the system needs about 50 primary cores plus about 50 redundant cores, or 100 always-on cores total. Running those cores for 24 hours/day at \$0.05/hr costs \$120/day.
3. **Batch requirements:** Process 3.6 TB (1-hour window) in 10 minutes. High throughput (sequential I/O) is efficient. The batch design needs 200 cores for 10 minutes per clock hour, accumulating 800 core-hours per day. At \$0.05/hr, that costs \$40/day.

Systems insight: Real-time is approximately 3× more expensive for the same data volume. This tax is justified only when the value of subsecond latency exceeds the added cost.

The ingestion paradigm also forces a choice on the model training side. Batch ingestion naturally pairs with periodic retraining: data accumulates in a store, a nightly or weekly job retrains the model on the updated dataset, and a new checkpoint is deployed. Stream ingestion makes that batch boundary less natural and raises the question of whether the model itself should update incrementally as each event arrives—a continuous learning approach where weights shift with the stream rather than resetting on a fixed schedule. Continuous weight updates from a live stream introduce risks that batch retraining avoids: a sudden distribution shift, an adversarial injection, or a burst of low-quality events can corrupt model weights before any validation gate intervenes. Microsoft’s public account of Tay illustrates the related danger of exposing a running system to untrusted live inputs, but it does not establish that Tay performed unrestricted online weight updates (Lee 2016). Most production systems therefore decouple the two concerns, streaming data into a buffer while still training on accumulated batches, reserving continuous online updates for models with narrow, well-monitored distributions.

Production systems must balance cost vs. latency trade-offs when selecting patterns: real-time processing is often materially more expensive than batch processing, commonly several times higher and sometimes an order of magnitude or more in total cost per byte processed. This cost differential arises from several factors: streaming systems require always-on infrastructure rather than schedulable resources; they maintain redundant processing for fault tolerance to ensure no events are lost; they need low-latency networking and storage to meet millisecond-scale SLAs; and they cannot benefit from the economies of scale that batch processing achieves by amortizing startup costs across large data volumes. A batch job processing one terabyte might use 100 machines for 10 minutes, while a streaming system processing the same data over 24 hours needs dedicated resources continuously available. This difference drives many architectural decisions about which data truly requires real-time processing. A cost estimate makes that premium explicit.

4.5.5.2 ETL and ELT comparison

Pipeline architecture dictates when and where computational resources are applied to raw incoming records. Compare the two sequence flows in figure 4.9 to see where the transformation stage executes in traditional ETL¹⁴ versus modern ELT¹⁵ pipelines.




Figure 4.9: ETL vs. ELT Architectural Paradigms: Shifting the transformation boundary relative to storage. ETL enforces schema validation and cleaning before loading to protect warehouse quality, whereas ELT loads raw data directly into scalable storage, trading warehouse compute and footprint for flexible, repeatable re-transformations.

For machine learning systems, ELT decouples data ingestion from feature engineering. Storing raw immutable events in the lakehouse allows data scientists to re-extract novel features over historical datasets without re-ingesting external sources, whereas rigid ETL transformations permanently discard unparsed fields that might later prove predictive.

The ELT pattern reverses this order, loading raw data first and applying transformations within the target system. For ML development, this enables flexible feature experimentation on the same raw data. Multiple teams can compute different aggregation windows, and when transformation logic bugs are discovered, teams reprocess by rerunning queries rather than re-ingesting from sources. This flexibility accelerates ML experimentation where feature engineering requirements evolve rapidly. The cost is higher storage requirements (raw data is larger than transformed data), repeated computation when multiple models transform the same source data, and greater complexity in enforcing privacy compliance when raw sensitive data persists in storage.

¹⁴ | **Extract, transform, load (ETL):** Conventional ETL validates *schema and data types*—deterministic checks that either pass or fail. ML ETL must additionally validate *distributional properties*: that the distribution of values has not shifted in a way that degrades model performance, requiring statistical tests (K-S test, PSI) rather than schema validation, where the pass/fail threshold is a business decision, not a technical one. The further trade-off is rigidity: changing feature computation logic requires reprocessing the entire dataset from raw sources, a cost that grows linearly with data volume.

¹⁵ | **Extract, load, transform (ELT):** This approach stores raw data before transformation, allowing teams to revise a transformation query and rerun it over the stored input without re-ingesting from source systems. The iteration time still depends on data volume, query complexity, warehouse capacity, and materialization strategy.

Napkin Math 4.6: The cost of transformation placement

Problem: A team processes 10 TB of raw clickstream data daily and must compute user session features for 3 ML models, each requiring different aggregation windows (1-hour, 24-hour, 7-day). Which is cheaper, ETL or ELT?

Math:

1. **ETL approach:** Transform before loading. Compute all three aggregation windows in a Spark cluster before loading into the warehouse.
 - Spark compute: 10 TB at \$5/TB = \$50/day
 - Storage: 3 transformed datasets, ~2 TB each = 6 TB at \$23/TB/month = \$138/month
 - Schema change cost: Re-run full pipeline (~4 hours) per change
2. **ELT approach:** Load raw data first, transform in warehouse.
 - Storage: 10 TB raw/day, 30 days retention = 300 TB at \$23/TB/month = \$6,900/month
 - Query compute: 3 models, each with \$5/query of daily query cost over 30 days, totals \$450/month
 - Schema change cost: Rewrite SQL query (~30 minutes) per change

Systems insight: ETL saves \$6,762/month in storage. After normalizing Spark compute to a monthly cost, ETL totals \$1,638/month (\$1,500/month compute + \$138/month storage), while ELT totals \$7,350/month (\$6,900/month storage + \$450/month query compute), for a cloud-cost advantage of \$5,712/month before engineering labor. ETL still costs 8× more engineering time per schema change. The exact break-even point depends on engineering labor cost and schema-change frequency: stable schemas favor ETL's lower cloud cost, while frequent feature-definition changes favor ELT's faster iteration.

Production ML systems rarely use one pattern exclusively. Structured data with stable schemas often flows through ETL for efficiency and compliance, while unstructured data or rapidly evolving feature pipelines benefit from ELT's flexibility. For deep learning workloads processing images, audio, or text, the ELT preference runs even deeper: the "Transform" step for unstructured data often executes inside the ML framework's data loader rather than in the warehouse at all, applying random crops, spectrogram generation, or text tokenization on-the-fly during each training epoch. Materializing ten augmented variants upstream would increase a 50,000-image dataset to 500,000 stored copies, so ELT's "load raw, transform late" principle extends naturally to the training loop itself. The accompanying ETL/ELT cost comparison makes the cost of transformation placement concrete.

When streaming components enter ETL/ELT architectures, the tool choice is really a failure-mode choice. The CAP theorem¹⁶ states that during a network partition, a distributed read/write service cannot guarantee both linearizable consistency and availability for every request. Apache Kafka¹⁷ provides ordering within each partition; its availability and durability during failures depend on acknowledgment, replication, and leader-election settings. Apache Pulsar separates stateless brokers from durable message storage in replicated Apache BookKeeper ledgers. Its built-in geo-replication is asynchronous, so a remote cluster may lag during failures or partitions rather than providing simultaneous strong cross-region consistency and availability. Amazon Kinesis exposes operational trade-offs through shard capacity, retention, and producer/consumer configuration, but under CAP it still *cannot* guarantee both strict consistency and availability during a network partition.

Choosing between upfront transformation and load-time transformation hinges on query reuse: transformed data that is queried hundreds of times amortizes ETL compute, whereas single-pass exploratory queries favor ELT warehouse elasticity. The parameterized cost model in listing 4.2 evaluates this break-even threshold across varying data volumes (V) and query frequencies (Q).

Listing 4.2: ETL vs. ELT Architectural Break-Even: Parameterized financial comparison computing compute and storage expenses across transformation strategies.

```
# ETL vs ELT cost comparison
daily_raw_tb = 10
s3_per_tb_mo = 23
spark_per_tb = 5
n_models = 3
retention_days = 30
query_cost_per = 5

# ETL
etl_spark_daily = daily_raw_tb * spark_per_tb
etl_datasets = 3
etl_tb_each = 2
etl_storage_tb = etl_datasets * etl_tb_each
etl_storage_mo = etl_storage_tb * s3_per_tb_mo

# ELT
elt_storage_tb = daily_raw_tb * retention_days
elt_storage_mo = elt_storage_tb * s3_per_tb_mo
elt_query_mo = n_models * query_cost_per * retention_days

# Savings
storage_savings_mo = elt_storage_mo - etl_storage_mo
```

¹⁶ | **CAP (consistency, availability, partition tolerance) theorem:** Conjectured by Brewer (2000) and formally proved by Gilbert and Lynch (2002). During a network partition, linearizable consistency may require rejecting requests, whereas an available system may return stale or divergent values. CAP does not guarantee feature-transformation parity or point-in-time correctness.

¹⁷ | **Apache Kafka:** Kafka uses a partitioned, leader-based log and orders records within each partition, not across a topic. Durability and write availability during failures depend on acknowledgment mode, replication, in-sync replica requirements, and leader-election behavior.

The storage difference favors ETL in this scenario, but the listing does not price schema-change labor; the worked calculation above treats that labor as a separate decision factor.

4.5.5.3 Feature computation placement

For ML pipelines, **feature computation placement** decides which resource pays for a feature: storage pays when features are materialized, while compute and latency pay when features are generated on demand. This choice significantly impacts training speed, storage costs, and reproducibility.

One approach is to precompute features during ETL and store the results. Pipeline-computed features offer fast training iteration (features are ready on disk), reproducibility (the same features are used consistently), and reduced training compute. The drawbacks are storage cost (features stored separately from raw data), staleness risk (precomputed features may diverge when logic changes), and inflexibility (any change requires full recomputation).

The alternative is computing features on the fly during training. Loader-computed features guarantee always-fresh computation (logic changes are immediately reflected), flexible experimentation (easy to modify features), and reduced storage (only raw data is stored). The cost is slower training (computation repeats each epoch), higher compute expenditure (GPUs often idle waiting for features), and potential nondeterminism if not carefully implemented.

In practice, hybrid patterns predominate. Expensive, stable features (user embeddings requiring matrix factorization, historical aggregations spanning months of data) are precomputed and materialized. Cheap, time-sensitive features such as recency signals, session context, and time-based transformations are computed in the data loader.

For example, a recommendation system precomputes stable user representation features (expensive, stable over days) while computing time-since-last-interaction features (cheap, time-sensitive) in the data loader. This balances storage costs, computation time, and feature freshness based on each feature's specific characteristics.

4.5.5.4 Integration strategies and KWS case study

Regardless of whether ETL or ELT approaches are used, integrating diverse data sources remains a core ingestion challenge. Data may originate from databases, APIs, file systems, and IoT devices, each with its own format (relational rows, JavaScript Object Notation (JSON)¹⁸ documents, binary streams), access protocol, and update frequency. The systems principle is to standardize at the ingestion boundary: normalize formats, validate schemas, and present a consistent interface to downstream processing regardless of source. This boundary standardization separates the complexity of source diversity from the complexity of feature engineering, allowing each to evolve independently.

KWS production systems often perform always-on wake-word detection on-device. Backend streaming and batch pipelines may ingest activated requests, consented diagnostics, crowdsourced recordings, synthetic data, and validated interactions for serving and training. Batch processing typically follows an ETL pattern where audio undergoes normalization, noise filtering, and segmentation into consistent durations before storage in training-optimized formats.

Error handling in voice interaction systems requires special attention. Dead letter queues store failed recognition attempts for subsequent analysis, revealing edge cases that need coverage in future model iterations. Each incoming audio sample must pass quality validation (signal-to-noise ratio, sample rate, duration bounds, speaker proximity) before entering the processing pipeline. Invalid samples route to analysis queues rather than being discarded, since these failures often indicate acoustic conditions underrepresented in training data. Valid samples flow through to real-time detection while simultaneously being logged for potential inclusion in future training data.

This ingestion architecture completes the boundary layer where external data enters our controlled pipeline. Ingested data, however reliably delivered, is still raw: audio at inconsistent sample rates, text with varying encodings, numeric features on incompatible scales. Transforming these heterogeneous records into a uniform, model-ready representation while guaranteeing that the exact same transformations apply during both training and serving is the next challenge.

4.6 Systematic Data Processing

The ingestion stage lexed raw streams into well-formed records; the next compiler phase is the *optimization pass*: systematic data processing. Just as compiler optimizations must preserve program

¹⁸ | **JSON (JavaScript object notation)**: The schema flexibility that makes JSON a common format for APIs creates a validation bottleneck at the ingestion boundary. Unlike binary formats with pre-defined schemas, every JSON document requires parsing and schema validation before it can be standardized for downstream use. This per-record overhead can make ingestion slower than with binary formats such as Protobuf; the difference depends on the schema, parser, payload, and workload.

semantics while improving performance, data transformations must preserve signal while improving model readiness. The governing constraint is consistency: transformations shared by training and serving must preserve equivalent semantics, repeated operations must be retry-safe, and processing must scale without losing lineage.

Sculley et al. (2015) identify data dependencies, changes in the external world, configuration debt, and missing monitoring as hidden technical debt risks in production ML systems. Training-serving inconsistency is one concrete form of this risk. Consider normalizing transaction amounts during training by removing currency symbols and converting to floats, but forgetting to apply identical preprocessing during serving. This seemingly minor inconsistency can materially degrade model accuracy. For our KWS system, the tension is concrete: transformations must standardize across diverse recording conditions (varying microphones, noise levels, sample rates) while preserving the acoustic characteristics that distinguish wake words from background speech—and this standardization must be identical in both training and serving paths.

4.6.1 Training-serving consistency

The training-serving consistency challenge extends beyond applying the same code—it requires that parameters computed on training data (normalization constants, encoding dictionaries, vocabulary mappings) are stored and reused during serving. This requirement is the **consistency imperative**.



Definition 4.3: The consistency imperative

The Consistency Imperative requires equivalent transformation behavior and state across training and serving environments.

1. **Significance:** KL divergence ($\mathcal{D}_{\text{KL}}(p_{g_{\text{serve}}} \| p_{g_{\text{train}}})$) between feature distributions induced by the serving transformation g_{serve} (current) and training transformation g_{train} (baseline) can signal misalignment. Its relationship to performance degradation depends on the model, task, and shifted features and must be checked against outcome evidence.
2. **Distinction:** Unlike data quality, which focuses on the cleanliness of a single record, the consistency imperative focuses on the alignment of the entire transformation pipeline.
3. **Common pitfall:** A frequent misconception is that consistency is “fixed” by sharing code. In reality, it is a state synchronization problem: parameters computed on training data (for example, means, standard deviations) must be stored and reused during serving.

The stakes are high: violating the consistency imperative silently degrades model accuracy in production. Before examining cleaning and transformation techniques, verify this requirement from the serving path backward.



Checkpoint 4.3: Defensive processing

Training-serving skew is a common cause of silent production degradation.

- ☐ **Definition:** can you define training-serving skew in terms of the training and live-request processing paths?
- ☐ **Mechanism:** can you explain why recomputing normalization means independently in training and serving can make live data look unfamiliar to the model?
- ☐ **Invariant:** can you explain how shared logic, synchronized state, and parity checks reduce skew?

Data cleaning is the first place where consistency can be either enforced or broken. Raw data frequently contains missing values, duplicates, or outliers that degrade model performance. The key insight is that cleaning operations must be deterministic and reproducible: given the same input, they must produce the same output regardless of environment. This requirement shapes which cleaning techniques are safe to use in production.

Data cleaning might involve removing duplicate records based on deterministic keys, handling missing values through imputation or deletion using rules that can be applied consistently, and correcting formatting inconsistencies systematically. For instance, a customer database might normalize “John Doe,” “john doe,” and “DOE, John” to a common display format. Normalization alone does not establish that records refer to the same person; identity resolution requires a stable identifier or a separately validated matching rule. Both the formatting and matching rules must be captured in code that executes equivalently in training and serving.

Outlier detection and treatment is another important aspect of data cleaning, but one that introduces consistency challenges. Outliers can sometimes represent valuable information about rare events, but they can also result from measurement errors or data corruption. ML practitioners must carefully consider the nature of their data and the requirements of their models when deciding how to handle outliers. Simple threshold-based outlier removal (removing values more than three standard deviations from the mean) maintains training-serving consistency if the mean and standard deviation are computed on training data and reused during serving. However, more sophisticated outlier detection methods that consider relationships between features or temporal patterns require careful engineering to ensure consistent application.

Quality assessment complements data cleaning by systematically evaluating the reliability and usefulness of data across multiple dimensions: accuracy, completeness, consistency, and timeliness. In production systems, data quality degrades in subtle ways that basic metrics miss: fields that never contain nulls suddenly show sparse patterns, numeric distributions drift from their training ranges, or categorical values appear that were not present during model development.

To address these subtle degradation patterns, production quality monitoring requires specific metrics beyond simple missing value counts (section 4.5.1). Critical indicators include null value patterns by feature (sudden increases suggest upstream failures), count anomalies ($10\times$ increases often indicate data duplication or pipeline errors), value range violations (prices becoming negative, ages exceeding realistic bounds), and join failure rates between data sources. Statistical drift detection¹⁹ becomes essential by monitoring means, variances, and quantiles of features over time to catch gradual degradation before it impacts model performance. For example, in an e-commerce recommendation system, the average user session length might gradually increase from eight minutes to 12 minutes over six months due to improved site design, but a sudden drop to three minutes suggests a data collection bug.

Tool sophistication matters less than whether quality standards remain identical across environments. Data profiling tools provide summary statistics and visualizations that help identify potential quality issues, while advanced techniques employ unsupervised learning algorithms to detect anomalies or inconsistencies in large datasets. The key is maintaining identical quality standards and validation logic across training and serving to prevent quality issues from creating training-serving skew.

Transformation techniques convert data from its raw form into a model-ready representation, but the systems risk lies in the parameters those transformations learn. Common transformation tasks include normalization and standardization, which scale numerical features to a common range or distribution. For example, square footage and room count may differ greatly in scale; normalization places them on comparable numerical ranges, though their influence still depends on the model and learned parameters (Bishop 2006). Maintaining training-serving consistency requires that normalization parameters computed on training data be stored and applied identically during serving. Operationally, these parameters must be persisted alongside the model itself and loaded during serving initialization.

Beyond numerical scaling, other transformations might involve encoding categorical variables, handling date and time data, or creating derived features. For instance, one-hot encoding is often used to convert categorical variables into a format that can be readily understood by many machine learning algorithms. Categorical encodings must handle both the categories present during training and unknown categories encountered during serving. A reliable approach computes the category vocabulary during training (the set of all observed categories), persists it with the model, and during serving either maps unknown categories to a special “unknown” token or uses default values. Without this discipline, serving encounters categories the model never saw during training, potentially causing errors or degraded performance.

¹⁹ | **Statistical drift detection:**

The means, variances, and quantiles tracked in quality monitoring are early-warning signals for the degradation equation’s divergence term $\mathcal{D}(P_t \| P_0)$. A mean session length shifting from eight to 12 minutes over six months may indicate drift; a sudden drop to three minutes may indicate a collection fault. Outcome checks and source investigation are needed before choosing retraining or a source-system fix.

A health prediction model may receive raw GPS coordinates for each patient visit, but latitude and longitude alone do not directly represent access to care. An engineer might derive *distance to nearest hospital* as a candidate proxy for geographic access. Whether that feature predicts outcomes, improves the model, or introduces sensitive-location and socioeconomic proxies must be established with data and domain review rather than assumed from the transformation alone.

Feature engineering is this act of using domain knowledge to create new features that make machine learning algorithms work more effectively (Kuhn and Johnson 2013). The step is often considered more art than science, requiring creativity and deep understanding of both the data and the problem at hand. Feature engineering might involve combining existing features, extracting information from complex data types, or creating entirely new features based on domain insights. In a retail recommendation system, for example, engineers might create features that capture the recency, frequency, and monetary (RFM) value of customer purchases, known as RFM analysis.

Feature engineering can be a high-leverage activity because it changes what signal the model can see. In some workloads, well-designed features improve performance more than another round of algorithm selection or hyperparameter tuning. That leverage must be balanced against the consistency requirements of production systems. Every engineered feature shared by training and serving must have equivalent behavior in both environments. Production systems therefore implement feature engineering logic in shared libraries or coordinated pipelines rather than reimplementing it independently. Many organizations use **feature stores**,²⁰ discussed in section 4.8.5, to reduce inconsistency across environments.

Applying these processing concepts to our KWS system: the audio recordings flowing through our ingestion pipeline, whether from crowdsourcing, synthetic generation, or real-world captures, require careful cleaning to ensure reliable wake word detection. Raw audio data often contains imperfections that our problem definition anticipated: background noise from various environments (quiet bedrooms to noisy industrial settings), clipped signals from recording level issues, varying volumes across different microphones and speakers, and inconsistent sampling rates from diverse capture devices. The cleaning pipeline must standardize these variations while preserving the acoustic characteristics that distinguish wake words from background speech, a quality-preservation requirement that directly impacts our 98 percent accuracy target.

Quality assessment for KWS extends the general principles with audio-specific metrics. Beyond checking for null values or schema conformance, our system tracks background noise levels (signal-to-noise ratio above 20 dB), audio clarity scores (frequency spectrum analysis), and speaking rate consistency (wake word duration within 500 ms–800 ms). The pipeline flags recordings whose noise, duration, clipping, or distortion falls outside the scenario’s acceptance rules. These screens do not prove that a sample is useful, but they reduce the chance that known defects reach model development. Recall how figure 4.1 demonstrated the compounding effects of early data quality failures; source checks reduce that risk before derived artifacts multiply it.

Transforming audio data for KWS involves converting raw waveforms into formats suitable for ML models while maintaining training-serving consistency. Raw audio waveforms (sequences of amplitude values sampled thousands of times per second) are high-dimensional. KWS pipelines often transform them into compact representations that emphasize the frequencies and temporal patterns most relevant to speech. Figure 4.10 compares three illustrative representations used in this pipeline: a waveform, its time-frequency spectrogram, and a compact MFCC approximation. These standardized feature representations, typically Mel-frequency cepstral coefficients (MFCCs)²¹ or spectrograms,²² emphasize speech-relevant characteristics while reducing noise and variability across different recording conditions.

4.6.2 Idempotent transformations

Reliability adds retry safety to these quality foundations. While quality focuses on what transformations produce, reliability also governs how safely they can be repeated. **Idempotency**²³ means that applying an operation repeatedly has the same effect as applying it once. This property proves essential for production ML systems where processing may be retried after failures, data may be reprocessed to fix bugs, or the same data may flow through multiple processing paths. Determinism is a related but distinct requirement: identical inputs and captured state should produce identical outputs.

²⁰ | **Feature store:** A core failure mode is training-serving skew: training features may be computed in batch (for example, seven-day rolling averages over historical data) while serving features are computed in real time from a streaming source; even with nominally shared logic, timing and state differences can create systematic discrepancies. A feature store can reduce this risk by coordinating definitions, data sources, timestamp handling, and offline and online materialization, but correctness still depends on implementation, freshness, and synchronization. Uber’s Michelangelo helped establish this dual-interface pattern for batch training and low-latency serving (Hermann and Del Balso 2017).

²¹ | **MFCC (mel-frequency cepstral coefficients):** This transformation achieves its compactness and noise resistance by applying mel-scale filtering, which selectively emphasizes the frequencies humans use to distinguish speech. This process reduces thousands of raw audio samples from a small time window (for example, 25 ms) into just 13–39 coefficients, the aggressive dimensionality reduction required for always-on, kilobyte-scale hardware. A parameter mismatch between training and serving creates feature skew and can degrade accuracy.

²² | **Spectrogram:** The short-time Fourier transform (STFT) computes this 2D representation by converting the one-dimensional waveform into a time-frequency image. This representation lets image-style ML models process audio, but it also creates a rigid dependency: a mismatch in STFT parameters (for example, a 25 ms vs. 30 ms window) changes the feature representation between training and serving and can degrade performance.




Figure 4.10: Audio Feature Representations: Three illustrative views show an audio waveform, its time-frequency spectrogram, and a compact MFCC approximation. The spectrogram encodes time, frequency, and signal power; the bottom panel arranges a smaller set of coefficients over time.

Consider a light switch to build intuition. Flipping the switch to the “on” position turns the light on. Flipping it to “on” again leaves the light on; the operation can be repeated without changing the outcome. This is idempotent behavior. In contrast, a toggle switch that changes state with each press is not idempotent: pressing it repeatedly alternates between on and off states. In data processing, we want light switch behavior where reapplying the same transformation yields the same result, not toggle switch behavior where repeated application changes the outcome unpredictably.

Idempotent transformations enable reliable error recovery. When a processing job fails midway, the system can safely retry processing the same data without worrying about duplicate transformations or inconsistent state. A nonidempotent transformation might append data to existing records, so retrying would create duplicates. An idempotent transformation would upsert data (insert if not exists, update if exists), so retrying produces the same final state. This distinction becomes critical in distributed systems where partial failures are common and retries are the primary recovery mechanism.

Handling partial processing failures requires careful state management. Processing pipelines should be designed so that each stage can be retried independently without affecting other stages. Checkpoint-restart mechanisms enable recovery from the last successful processing state rather than restarting from scratch. For long-running data processing jobs operating on terabyte-scale datasets, checkpointing progress every few minutes means a failure near the end requires reprocessing only recent data rather than the entire dataset. The checkpoint logic must carefully track what data has been processed and what remains, ensuring no data is lost or processed twice.

Deterministic transformations are those that always produce the same output for the same input, without dependence on external factors like time, random numbers, or mutable global state. Transformations that depend on current time (for example, computing “days since event” based on current

23 | Idempotency: From Latin *idem* (“the same”) + *potens* (“having power”) – literally, “having the same power when applied again.” In ML pipelines, idempotency enables safe retries after partial failures: a nonidempotent operation (for example, appending to a log) creates duplicates on retry. An idempotent operation leaves the same resulting state after one application or repeated applications.

date) break determinism because reprocessing historical data would produce different results. The solution is to capture temporal reference points explicitly: instead of “days since event,” compute “days from event to reference date” where reference date is fixed and persisted. Random operations should use seeded random number generators where the seed is derived deterministically from input data, ensuring reproducibility.

Reliability in the KWS pipeline requires reproducible feature extraction. Audio preprocessing must be deterministic: given the same raw audio file, the same MFCC features are always computed regardless of when processing occurs or which server executes it. This enables debugging model behavior (can always recreate exact features for a problematic example), reprocessing data when bugs are fixed (produces consistent results), and distributed processing (different workers produce identical features from the same input). The processing code captures all parameters (fast Fourier transform (FFT) window size, hop length, number of MFCC coefficients) in configuration versioned alongside the code, ensuring reproducibility across time and execution environments. Even with rigorous design, production systems must implement runtime monitoring to detect skew if it emerges; chapter 14 covers operational comparison and distribution monitoring at scale.

4.6.3 Distributed processing

Scale becomes the next constraint once quality and reliability are in place. Quality ensures transformations produce *correct* outputs; reliability ensures they produce *consistent* outputs. Neither matters if processing cannot keep pace with data volume or its deadline. As datasets and experiment concurrency grow, data processing can outgrow one machine even when the raw bytes still fit on local storage. The cleaning techniques that work on gigabytes in memory must then become partition-aware, out-of-core, or distributed without changing their semantics.


These challenges manifest when quality assessment must keep pace with incoming data, when feature engineering requires computing statistics across entire datasets before transforming individual records, and when transformation pipelines create bottlenecks at massive volumes. Processing must scale from development (gigabytes on laptops) through production (terabytes across clusters) while maintaining consistent behavior.

To address these scaling bottlenecks, data must be partitioned across multiple computing resources, which introduces coordination challenges. Distributed coordination is constrained by network round-trip times: local operations complete in microseconds while network coordination requires milliseconds, creating a $1,000\times$ latency difference. This constraint explains why operations requiring global coordination (like computing normalization statistics across 100 machines) create bottlenecks. Each partition computes local statistics quickly, but combining them requires information from all partitions.

Data locality becomes critical at this scale. At 10 GB/s peak throughput, transferring one terabyte of training data across a network takes on the order of 100 seconds; reading the same amount from a 5 GB/s SSD takes on the order of 200 seconds. These are the same order of magnitude, which drives ML system design toward compute-follows-data architectures.²⁴ When processing nodes access local data at RAM speeds (50–200 GB/s) but must coordinate over networks limited to 1–10 GB/s, the bandwidth mismatch creates severe bottlenecks. Geographic distribution amplifies these challenges: cross-data center coordination must handle network latency (50–200 ms between regions), partial failures, and regulatory constraints preventing data from crossing borders. Understanding which operations parallelize easily vs. those requiring expensive coordination determines system architecture and performance characteristics. This overhead constitutes a **coordination tax** that limits distributed data processing. The size of this tax, and whether centralizing or aggregating is faster, depends on the ratio between network round-trip and local compute time—derived in napkin math 4.7.

Single-machine processing suffices for surprisingly large workloads when engineered carefully. Modern servers with 256 gigabytes RAM can process datasets of several terabytes using out-of-core processing that streams data from disk. Libraries like Dask or Vaex enable pandas-like APIs that automatically stream and parallelize computations across multiple cores. Before investing in distributed processing infrastructure, teams should exhaust single-machine optimization: using efficient data formats (Parquet²⁵ instead of CSV), minimizing memory allocations, using vectorized operations, and exploiting multi-core parallelism. The operational simplicity of single-machine

²⁴ | **MapReduce:** Designed by Dean and Ghemawat (2004) at Google to process the company’s multi-petabyte web index across thousands of commodity machines. Its scheduler preferentially assigns map tasks to nodes that hold the input data, reducing input-network traffic. This locality-aware pattern influenced Hadoop and later distributed data-processing systems.



Local aggregation beats gather-all when network latency dominates.

²⁵ | **Parquet:** A columnar storage format that organizes data by column, not by row. For the single-machine optimizations described, this is critical; instead of wastefully reading an entire CSV row to access a few columns, a Parquet reader loads only the specific columns needed for a computation. The resulting I/O reduction depends on the selected columns, encoding, compression, and workload.

processing (no network coordination, no partial failures, simple debugging) makes it preferable when performance is adequate.

Napkin Math 4.7: The coordination tax

Problem: Mean normalization must be computed across 1 TB of features distributed across 100 nodes. Is it faster to (A) gather all data to one node and compute centrally, or (B) compute local means and aggregate them?

Option A (centralized):

- Transfer 1 TB at 10 GB/s network: 100 seconds
- Compute mean on single node: ~5–20 seconds at RAM bandwidth
- Total: ~105–120 seconds

Option B (distributed):

- Each node computes local mean: ~0.05–0.2 seconds (10 GB at RAM speed)
- Send 100 partial means (8 bytes each): < 1 ms
- Aggregate: negligible
- Total: ~0.05–0.2 seconds (hundreds to a few thousand times faster)

Systems insight: Algebraically decomposable reductions can often use partial aggregation near the data; for example, a distributed mean carries partial sums and counts. Joins and cross-products may require shuffling, although partitioning and colocation can reduce that cost. Pipeline design should minimize unnecessary movement by pushing suitable computation toward the data, the compute-follows-data principle central to systems like MapReduce (Dean and Ghemawat 2004), Spark (Zaharia et al. 2010), and modern ML frameworks.

Parallelizing data ingestion across distributed workers yields diminishing throughput returns whenever unpartitionable bottlenecks—such as sequential decompression or central database lock acquisition—limit concurrency. Amdahl’s Law²⁶ formalizes this scaling ceiling in equation 4.4:

$$\text{Speedup} \leq \frac{1}{f_{\text{serial}} + \frac{f_{\text{parallel}}}{N_{\text{workers}}}} \quad (4.4)$$

where f_{serial} denotes the execution fraction spent in inherently sequential stages, f_{parallel} represents the parallelizable fraction ($f_{\text{serial}} + f_{\text{parallel}} = 1$), and N_{workers} denotes active worker processes. Even when 90 percent of a feature extraction pipeline parallelizes across hundreds of nodes ($f_{\text{serial}} = 0.10$), the maximum theoretical speedup cannot exceed 10×, demonstrating why eliminating serial I/O bottlenecks takes precedence over adding compute nodes.

Framework choice follows the same coordination question as the Amdahl analysis: it depends on which parts of the transformation can run independently and which parts need shared state or ordering. Apache Spark parallelizes transformations across clusters of machines, handling data partitioning, task scheduling, and fault tolerance automatically. Beam provides a unified API for both batch and streaming processing, enabling the same transformation logic to run on multiple execution engines (Spark, Flink, Dataflow). TensorFlow’s tf.data API optimizes data loading pipelines for ML training, supporting distributed reading, prefetching, and transformation. The choice depends on whether processing is batch or streaming, how transformations parallelize, and what execution environment is available.

The feature computation placement trade-off introduced in section 4.5.5.3 takes on additional significance at scale. When distributed processing increases throughput, the cost of recomputing features across hundreds of workers per epoch must be weighed against the storage cost of materializing those features once. At terabyte scale, even small per-example compute costs multiply into significant overhead, reinforcing why production systems adopt hybrid patterns: precomputing expensive, stable features while computing cheap, time-sensitive features on-the-fly.

Scalability in the KWS pipeline manifests at multiple stages. Development uses single-machine processing on sample datasets to iterate rapidly. Training at scale may require distributed processing

²⁶ | **Amdahl’s law:** Amdahl (1967) presented the serial-fraction argument at the 1967 AFIPS Spring Joint Computer Conference to explain why multiprocessor designs face diminishing returns. The original point – that the serial fraction of a workload imposes a hard ceiling on parallelism – applies directly to data pipelines: operations like computing global normalization statistics force serial aggregation phases that cap speedup regardless of how many workers process individual records in parallel.

when the dataset (23.4 million examples) exceeds a machine's practical capacity or when experiments run concurrently. Per-file feature extraction is embarrassingly parallel, but end-to-end speedup is bounded by shared storage bandwidth, metadata operations, scheduling, and output contention. Production deployment adds a stricter 16 KB preprocessing-state budget alongside the 64 KB model-size limit, necessitating careful footprint optimization to fit processing within device capabilities.

4.6.4 Transformation lineage

Governance completes the four-pillar view of data processing by ensuring accountability and reproducibility. The governance pillar requires tracking what transformations were applied, when they executed, which version of processing code ran, and what parameters were used. This **transformation lineage**²⁷ supports reproducibility, debugging, auditability, and iterative improvement when transformation bugs are discovered. It can supply evidence for documentation or compliance workflows, but lineage alone does not explain a model's decision or establish regulatory compliance.

Transformation versioning captures which version of processing code produced each dataset. When transformation logic changes (fixing a bug, adding features, or improving quality), the version number increments. Datasets are tagged with the transformation version that created them, enabling identification of all data requiring reprocessing when bugs are fixed. This versioning extends beyond just code versions to capture the entire processing environment: library versions (different NumPy versions may produce slightly different numerical results), runtime configurations (environment variables affecting behavior), and execution infrastructure (CPU architecture affecting floating-point precision).

Parameter tracking maintains the specific values used during transformation. For normalization, this means storing the mean and standard deviation computed on training data. For categorical encoding, this means storing the vocabulary (set of all observed categories). For feature engineering, this means storing any constants, thresholds, or parameters used in feature computation. These parameters are typically serialized alongside model artifacts, ensuring serving uses identical parameters to training. Modern ML frameworks like TensorFlow and PyTorch provide mechanisms for bundling preprocessing parameters with models, simplifying deployment and ensuring consistency.

Processing lineage for reproducibility tracks the complete transformation history from raw data to final features. This includes which raw data files were read, what transformations were applied in what order, what parameters were used, and when processing occurred. Lineage systems like Apache Atlas, Amundsen, or commercial offerings instrument pipelines to automatically capture this flow. When model predictions prove incorrect, engineers can trace back through lineage to identify the training data that contributed to the behavior, the quality scores attached to that data, the transformations applied, and whether the exact scenario can be recreated for investigation.

Code version ties processing results to the exact code that produced them. When processing code lives in version control (Git), each dataset should record the commit hash of the code that created it. This enables recreating the exact processing environment: checking out the specific code version, installing dependencies listed at that version, and running processing with identical parameters. Container technologies like Docker simplify this by capturing the entire processing environment (code, dependencies, system libraries) in an immutable image that can be rerun months or years later with identical results.

The governance pillar in the KWS pipeline tracks audio processing parameters that critically affect model behavior. When audio is normalized to standard volume, the reference volume level is persisted. When FFT transforms audio to frequency domain, the window size, hop length, and window function (Hamming, Hanning, etc.) are recorded. When MFCCs are computed, the number of coefficients, frequency range, and mel filterbank parameters are captured. This comprehensive parameter tracking enables several critical capabilities: reproducing training data exactly when debugging model failures, validating that serving uses identical preprocessing to training, and systematically studying how preprocessing choices affect model accuracy. Without this governance infrastructure, teams resort to manual documentation that inevitably becomes outdated or incorrect, leading to subtle training-serving skew that degrades production performance.

Clean, normalized, feature-ready data is still inert without meaning. That boundary matters operationally because a transformation can be perfectly reproducible yet attach the wrong target, and pipeline determinism cannot repair a mislabeled example. The remaining question is how to

²⁷ | **Data lineage:** When a model produces an erroneous or discriminatory prediction, lineage helps engineers determine whether the contributing problem originated in raw data, feature computation, training, or serving. The resulting trace can support incident investigation and jurisdiction-specific documentation duties, but its usefulness depends on complete instrumentation and retained history; it does not guarantee an explanation of model behavior.

assign meaning: labels declare which audio clips contain the wake word and which are background noise, introducing human judgment into what has been an automated pipeline.

4.7 Data Labeling

The processing pipelines in section 4.6 transform raw data into structured features, but supervised learning still requires labels that tell the model which patterns correspond to each target. Consider our KWS system: the ingestion and processing stages have produced millions of clean, standardized audio spectrograms, but training requires someone, or something, to declare which contain the wake word and which are background noise. This declaration is the **ground truth**,²⁸ and producing it at scale is often the most human-dependent and error-prone stage of the pipeline.

Unlike automated transformations that can be parallelized across machines, labeling introduces human judgment into the pipeline, creating unique engineering challenges. A crowdsourced annotator might mislabel a whispered “Alexa” as background noise. An expert radiologist might disagree with a colleague about a borderline diagnosis. Such disagreement can reflect genuine ambiguity, unclear instructions, annotator error, or differences in expertise; the labeling system must distinguish and manage these causes rather than treating every disagreement as irreducible. The infrastructure must therefore handle throughput, quality control, cost management, and governance.

4.7.1 Label types and system requirements

Building effective labeling systems requires understanding how different label types affect system architecture and resource requirements. Consider a practical example: building a smart city system that needs to detect and track various objects like vehicles, pedestrians, and traffic signs from video feeds. Labels capture information about key tasks or concepts, with each label type imposing distinct storage, computation, and validation requirements.

Classification labels represent the simplest form, categorizing images with a specific tag or (in multi-label classification) tags such as labeling an image as “car” or “pedestrian.” While conceptually straightforward, a production system processing millions of video frames must efficiently store and retrieve these labels. Storage requirements are modest (a single integer or string per image), but retrieval patterns matter: training often samples random subsets while validation requires sequential access to all labels, driving different indexing strategies.


Bounding boxes extend beyond simple classification by identifying object locations, drawing a box around each object of interest. Our system must track both object identity and spatial location within each frame. This spatial information introduces new storage and processing challenges, especially when tracking moving objects across video frames. Each bounding box stores four coordinates (x, y, width, height) plus the object class, so storage scales with the number of annotated objects per image rather than a fixed-size label per image. Bounding box annotation requires pixel-precise positioning that takes 10–20× longer than classification, dramatically affecting labeling throughput and cost.

Segmentation maps provide the most comprehensive information by classifying objects at the pixel level, highlighting each object in a distinct color. For our traffic monitoring system, this might mean precisely outlining each vehicle, pedestrian, and road sign. These detailed annotations significantly increase our storage and processing requirements. A segmentation mask for a 1920×1080 image requires about 2.1M labels (one per pixel), compared to perhaps 10 bounding boxes or a single classification label. If each box stores 4 coordinates, that is roughly 51,840× more scalar label entries than 10 boxes before accounting for per-value encoding, and the hours required per image for manual segmentation make this approach suitable only when pixel-level precision is essential.

Figure 4.11 contrasts five annotation modalities, and the choice depends on the input, task, and available annotation budget. Classification can label an entire scene, while bounding boxes and segmentation maps localize objects or regions. Production datasets may combine modalities: a single camera frame might carry a scene label, obstacle boxes, and path-region masks, with each label type serving a distinct downstream task.

Beyond these geometric labels, production systems must also manage rich metadata essential for quality control and debugging. The Common Voice dataset (Ardila et al. 2020) exemplifies this in speech recognition: tracking speaker demographics for fairness, recording quality metrics for filtering, and language information for multilingual support. If our traffic monitoring system fails in rainy conditions, weather metadata captured during collection pinpoints the coverage gap.

²⁸ | **Ground truth:** From remote sensing, where orbital measurements are verified by sending a team to the physical location – the “ground” – to establish the “truth.” The etymology carries a systems warning: ML labels are proxies for reality, not reality itself. When a crowdsourced annotator labels an image as “cat,” that label reflects the annotator’s judgment, not an objective fact. Every downstream metric – accuracy, precision, recall – is measured against this proxy, meaning label quality errors propagate silently into every evaluation of the model.



Finer annotation granularity multiplies storage and processing scale.

Label Type	Input Type	Output Type
Classification Label		Dog, Blanket, No cat
Bounding Box		
Segmentation Map		
Caption		A dog curls up on a spotted purple blanket.
Transcript		Once upon a time. . .

Figure 4.11: Data Annotation Modalities: Four rows show classification, bounding-box, segmentation, and caption labels for one image; the fifth shows a transcript label for an audio input. The appropriate annotation representation depends on the input modality and task.

This metadata requirement demonstrates how label type choice cascades through entire system design: the infrastructure must optimize storage for the chosen format, implement appropriate retrieval patterns, and track which model versions used which label versions to correlate quality improvements with performance gains.

4.7.2 Label accuracy and consensus

In the labeling domain, label quality centers on ensuring label accuracy despite the inherent subjectivity and ambiguity in many labeling tasks. Even with clear guidelines and careful system design, some fraction of labels will inevitably be incorrect (Northcutt et al. 2021; Thyagarajan et al. 2022). The challenge is not eliminating labeling errors entirely (an impossible goal) but systematically measuring inter-annotator agreement and managing error rates to keep them within bounds that do not degrade model performance.

Labeling failures arise from two distinct sources requiring different engineering responses. Figure 4.12 presents concrete examples of both failure modes. Some examples reflect degraded or ambiguous inputs where the correct label is difficult to infer from the data alone; others are visually clear but require domain knowledge, dataset-specific semantics, or expert judgment to label correctly. These different failure modes drive architectural decisions about annotator qualification, task routing, and consensus mechanisms: quality-based errors call for upstream data filtering, while expertise-based errors call for tiered annotator routing.



Figure 4.12: Pervasive Label Errors: Confidently mislabeled examples from widely used benchmarks, each showing the given label against the corrected ground truth (MNIST, CIFAR-10, CIFAR-100, Caltech-256, ImageNet, QuickDraw). Fine-grained confusions, such as a crab labeled lobster or a black stork labeled white stork, show that even curated datasets carry systematic errors, motivating careful quality control and expert annotation. Source: (Northcutt et al. 2021).

Given these inherent quality challenges, production ML systems implement multiple layers of quality control. Systematic quality checks continuously monitor the labeling pipeline through random sampling of labeled data for expert review and statistical methods to flag potential errors. The infrastructure must efficiently process these checks across millions of examples without creating bottlenecks. Sampling strategies typically validate 1-10 percent of labels, balancing detection sensitivity

against review costs. Higher-risk applications like medical diagnosis or autonomous vehicles may validate 100 percent of labels through multiple independent reviews, while lower-stakes applications like product recommendations may validate only 1 percent through spot checks.

Beyond random sampling approaches, collecting multiple labels per data point, often referred to as “consensus labeling,” can help identify controversial or ambiguous cases. Commercial labeling platforms such as Labelbox and Scale AI expose consensus and tiered quality-control workflows (Labelbox, Inc. 2024; Scale AI, Inc. 2024), but the statistical core is inter-annotator agreement. The consensus infrastructure typically collects several labels per example, computing metrics like Fleiss’ kappa, a generalization of the Cohen’s kappa statistic introduced in section 4.5.3 from two raters to any number of annotators (Fleiss 1971). Examples with low agreement, using thresholds such as the Landis-Koch bands as operational heuristics, route to expert review rather than forcing consensus from genuinely ambiguous cases (Landis and Koch 1977).

The consensus approach reflects an economic trade-off essential for scalable systems. Expert review costs more per example than crowdsourced labeling, but forcing agreement on ambiguous examples through majority voting of nonexperts can produce systematically biased labels. By routing only genuinely ambiguous cases to experts, identified through low inter-annotator agreement or failed gold-standard checks, systems balance cost against quality. This tiered approach enables processing millions of examples economically while maintaining quality standards through targeted expert intervention.

While technical infrastructure provides the foundation for quality control, successful labeling systems must also consider human factors. When working with annotators, organizations need reliable systems for training and guidance. This includes good documentation with clear examples of correct labeling, visual demonstrations of edge cases and how to handle them, regular feedback mechanisms showing annotators their accuracy on gold standard examples, and calibration sessions where annotators discuss ambiguous cases to develop shared understanding. For complex or domain-specific tasks, the system might implement tiered access levels, routing challenging cases to annotators with appropriate expertise based on their demonstrated accuracy on similar examples.

Quality monitoring generates substantial data that must be efficiently processed and tracked. The most informative signals span several dimensions. Inter-annotator agreement rates reveal whether multiple annotators converge on the same example, while label confidence scores capture how certain annotators feel about their decisions. Time per annotation serves as a dual-sided indicator: annotations completed too quickly suggest carelessness, while those taking too long suggest confusion or unclear guidelines. Error patterns expose systematic biases or misunderstandings in the annotator pool, and annotator performance on gold standard examples provides ground-truth calibration. Finally, demographic analysis of annotator behavior detects whether certain groups systematically label differently, which could introduce unintended bias into the training data. These metrics must be computed and updated efficiently across millions of examples, often requiring dedicated analytics pipelines that process labeling data in near real-time to catch quality issues before they affect large volumes of data.

4.7.3 Scaling with AI-assisted labeling

The scalability pillar drives AI assistance as a force multiplier for human labeling rather than a replacement. Manual annotation alone may not keep pace with a system’s data needs, while fully automated labeling cannot supply reliable judgment for every ambiguous or high-stakes case. AI-assisted labeling occupies the space between these extremes: using automation to handle clear cases and accelerate annotation while preserving human review where it matters. Figure 4.13 maps four common paths. Traditional supervision uses direct human labels. Semi-supervised learning exploits structure in unlabeled data alongside a labeled subset. Weak supervision replaces some individual annotations with programmatic labeling functions. Transfer learning reuses representations learned on another task, often reducing but not eliminating task-specific labels. Each path changes where supervision cost is paid and which assumptions, unlabeled-data structure, or pretrained artifacts the pipeline must validate.

These paths exchange different assumptions, not just label counts. Depending on the path, the pipeline needs task-relevant structure in unlabeled data, a transferable representation, measurable labeling-function errors, or an affordable scoring loop. Production systems should record each




Figure 4.13: AI-Augmented Labeling Decision Hierarchy: A top-level question about obtaining labeled data branches into four paths: traditional supervision, semi-supervised learning, weak supervision, and transfer learning, with active learning as a cost-saving alternative. The branches distinguish where supervision comes from and what additional assumptions or pretrained artifacts the pipeline must validate. Adapted from the Stanford AI Lab weak-supervision overview (Ratner et al. 2019).

assumption and test it against held-out, human-reviewed data. The diagram is a decision map, not a quality ranking.

AI-assisted labeling divides work according to the strengths of each participant. Humans judge ambiguous cases, catch subtle errors, and apply domain knowledge; models can process routine cases at scale, subject to validation. Production systems combine these capabilities through several complementary approaches.

Pre-annotation uses AI models to generate preliminary labels that humans then review and correct – transforming the task from “label from scratch” to “verify and fix.” Programmatic labeling frameworks like Snorkel (Ratner et al. 2018; Ratner et al. 2017) extend this further through weak supervision,²⁹ automatically generating initial labels at scale through rule-based heuristics, knowledge bases, and existing model outputs. In autonomous driving, pretrained object detection models can label vehicles and pedestrians that human annotators verify and refine, handling many clear cases automatically.

Large language models (LLMs) now assist labeling pipelines by generating descriptions, drafting labeling guidelines from examples, and explaining their reasoning for label assignments. Content moderation systems, for instance, use LLMs for initial content classification with explanations that human reviewers validate. However, LLM integration introduces systems challenges: provider- and model-dependent inference costs and rate limits, as well as the need for systematic output validation because LLMs can produce confident but incorrect labels. Many organizations adopt tiered approaches, using smaller specialized models for routine cases while reserving larger LLMs for complex scenarios requiring nuanced judgment.

Active learning makes the complementary trade-off: it spends model inference to reduce human labeling, so the relevant question is whether that compute fits inside the same budget.

Methods such as active learning³⁰ complement these approaches by intelligently prioritizing which examples need human attention (Settles 2009; Coleman et al. 2022). These systems continuously analyze model uncertainty to identify valuable labeling candidates. Rather than labeling a random sample of unlabeled data, active learning selects examples where the current model is most uncertain or where labels would most improve model performance. The infrastructure must efficiently compute uncertainty metrics (often prediction entropy or disagreement between ensemble models), maintain task queues ordered by informativeness, and adapt prioritization strategies based on incoming labels. Consider a medical imaging system: active learning might identify unusual pathologies for expert review while handling routine cases through preannotation that experts merely verify. This approach can substantially reduce required annotations in favorable settings, though it requires careful engineering to prevent feedback loops where the model’s uncertainty biases which data gets labeled. A budget calculation makes that leverage concrete.

²⁹ | **Weak supervision:** A “data programming” paradigm that exchanges some manual annotation labor for the upfront effort of writing and validating programmatic labeling functions. Once written, a function can be applied cheaply at scale, but execution, conflict resolution, quality monitoring, and maintenance still incur cost; the useful comparison is workload-specific rather than literally zero marginal cost.

³⁰ | **Active learning:** Inverts the traditional labeling paradigm: instead of randomly selecting examples to label, the model queries for the examples it needs most, typically those where prediction uncertainty is highest (Settles 2009). This can reduce the number of labels needed to reach a target accuracy, but the infrastructure trade-off is compute for labels: at \$0.01/image, scoring a full 10M pool costs \$100K before any human labels. Active learning becomes budget leverage only when the candidate pool is pre-filtered, inference is much cheaper, or the compute budget is separate from the labeling budget.

Napkin Math 4.8: The active learning multiplier

Problem: A 10M dataset has a \$50K labeling budget. Random sampling achieves 85 percent accuracy with 100K, while the target is 95 percent accuracy.

Physics:

1. **Sample efficiency:** Active learning can achieve target accuracy with fewer samples than random selection in favorable settings.
2. **Cost per point:** Random sampling = \$0.50/label. Active learning adds compute cost (~\$0.01/image for inference) to find hard examples.
3. **Multiplier:**
 - **Random:** Reaching 95 percent may require 1M labels (\$500K). Budget exceeded.
 - **Active labels only:** The system may need ~100K–200K hard examples (\$50K–\$100K).
 - **Full-pool scoring:** Scoring all 10M candidate images adds \$100K, so total active-learning cost becomes \$150K–\$200K. Budget exceeded.

Systems insight: Algorithm choice is a major lever on label count, but compute must be inside the budget model. Spending 10 percent of this budget on inference (\$5K) scores only 500K candidate images at the stated inference price and leaves room for about 90K labels. Active learning is viable only if the candidate pool is narrowed before scoring, inference cost drops substantially, or compute is funded separately from labeling.


Quality control becomes increasingly important as these AI components interact. The system must monitor both AI and human performance through systematic metrics. Model confidence calibration matters: if the AI reports 95 percent confidence but achieves only 75 percent accuracy at that confidence level, preannotations mislead human reviewers. Human-AI agreement rates reveal whether AI assistance helps or hinders: when humans frequently override AI suggestions, the preannotations may be introducing bias rather than accelerating work. These metrics require careful instrumentation throughout the labeling pipeline, tracking both final labels and the interaction between human annotators and AI at each stage.

These principles manifest at scale across safety-critical domains. Autonomous vehicle labeling infrastructure can process large volumes of sensor frames, using AI preannotation to label common objects while routing unusual scenarios (construction zones, emergency vehicles) to human experts—a distributed architecture where preannotation runs on GPU clusters while human review scales across annotation teams. Medical imaging systems face a parallel label-scarcity problem: large repositories of unlabeled clinical data make expert annotation the bottleneck and motivate data-efficient workflows that learn useful representations from unlabeled structure before expert labels are applied (Krishnan et al. 2022). Across such domains, the common data-engineering pattern is *tiered escalation*: automation handles clear cases, humans handle ambiguous ones, and monitoring ensures the boundary between “clear” and “ambiguous” adapts as both AI capability and deployment conditions evolve.

4.7.4 Automated labeling in KWS

Labeling our KWS data at scale creates a speech-specific challenge. Generating millions of labeled wake word samples without proportional human annotation cost requires moving beyond the manual and crowdsourced approaches introduced earlier. The Multilingual Spoken Words Corpus (MSWC) (Mazumder et al. 2021) demonstrates how automated labeling addresses this challenge through its innovative approach to generating labeled wake word data; the corpus contains over 23.4 million examples of one-second speech across 340,000 keywords in 50 languages.

This scale makes manual annotation impractical: 23.4 million examples at even 10 seconds per label would require approximately 65,000 hours, roughly 32.5 person-years of full-time effort. Broad, documented sourcing across 50 languages can improve coverage, but language count alone does not guarantee representative speakers, accents, devices, or acoustic environments. The automated system in figure 4.14 addresses the scale problem by starting with paired sentence audio and transcriptions, then using forced alignment³¹ to estimate word boundaries within continuous speech.



Full-pool scoring can exceed the label budget before labeling begins.

³¹ | **Forced alignment:** Given a known transcription, an aligner estimates where its words occur in the audio, often using frame-level acoustic scores and sequence decoding. Timing resolution and accuracy depend on the model, features, language, and recording quality, so boundary estimates still need quality checks. The known transcript makes word-level segmentation far cheaper than labeling every clip manually, though computation, review, and error correction remain.

At this scale, quality control becomes a sampling and provenance problem rather than disappearing. The pipeline should report coverage by language, speaker, accent, device, and acoustic environment; retain the source recording, transcript version, aligner version, and boundary confidence for each extracted clip; and route low-confidence alignments and underrepresented slices to human review. Held-out evaluation must then measure the intended wake words across those slices. These controls distinguish a large corpus from a representative one and make extraction errors traceable when a downstream KWS model fails.




Figure 4.14: Multilingual Data Preparation: Forced alignment and segmentation transform paired audio-text data into labeled one-second segments, creating a large-scale corpus for training keyword spotting models across 50 languages. This automated process enables scalable development of KWS systems by efficiently generating training examples from readily available speech resources like Common Voice and multilingual captioned content. Source: (Mazumder 2021).

The extraction system uses these precise timing markers to generate clean keyword samples while handling the engineering challenges our problem definition anticipated: background noise interfering with word boundaries, speakers stretching or compressing words unexpectedly beyond our target 500 ms–800 ms duration, and longer words exceeding the one-second boundary. MSWC provides automated quality assessment that analyzes audio characteristics to identify potential issues with recording quality, speech clarity, or background noise, which is essential for maintaining consistent standards across 23.4 million samples without the manual review expenses that would make this scale prohibitive.

Modern voice assistant developers often build on this automated labeling foundation. While automated corpora may not contain the specific wake words a product requires, they provide starting points for KWS prototyping, particularly in underserved languages where commercial datasets do not exist. Production systems typically layer targeted human recording and verification for challenging cases (unusual accents, rare words, or difficult acoustic environments), coordinating between automated processing and human expertise.

The pipeline has now produced its compilation artifacts: millions of feature vectors paired with ground truth labels. The question shifts from *what* data has been collected to *where* it lives and *how fast* it reaches the accelerators. *Storage architecture determines whether expensive GPUs spend their time computing or waiting.*

4.8 Storage Architecture

The labeled datasets from our pipeline (23.4 million samples spanning 50 languages for KWS) now require strategic storage decisions that determine training efficiency, serving latency, and long-term maintainability. Storage architecture addresses a core tension: batch training requires sequential scans across millions of examples, while real-time serving demands millisecond lookups of individual feature vectors. These competing access patterns shape every storage decision.

ML storage requirements diverge from those of transactional systems. Rather than optimizing for frequent small writes and point lookups that characterize e-commerce or banking, ML workloads

prioritize high-throughput sequential reads, large-scale scans, and schema flexibility. A database serving an e-commerce application performs well with millions of individual product lookups per second, but an ML training job scanning that entire catalog repeatedly across epochs requires completely different storage optimization.

4.8.1 Storage system options

Batch training scans millions of examples sequentially; real-time serving fetches one feature vector at a time. These opposing access patterns pull storage in two directions, and selecting a system means minimizing the data term ($\frac{D_{\text{vol}}}{\text{BW}}$) of the iron law of ML systems for whichever pattern dominates. Every storage medium imposes physical constraints on bandwidth that determine the maximum speed of the training and serving pipelines.

Two storage performance metrics govern this optimization. **IOPS (input/output operations per second)** counts the distinct read/write requests a device can handle per second, so it limits *random access* workloads such as fetching small batches of images or individual user profiles. **Throughput (bandwidth)** measures the volume of data transferred per second, typically $\text{IOPS} \times \text{Block Size}$, so it limits *sequential access* workloads such as scanning a Parquet file for training.

The choice between databases, data warehouses, and data lakes is fundamentally a choice about which of these metrics to optimize. Databases (online transaction processing [OLTP] systems) optimize for high IOPS with small block sizes, making them suited for serving individual feature vectors in real-time where per-request latency dominates. Data warehouses (online analytical processing [OLAP] systems) optimize for high throughput with large block sizes and sequential access, making them ideal for feature engineering and batch analytics. Data lakes prioritize capacity and throughput for unstructured data, essential for training jobs where the D_{vol} numerator is measured in petabytes and aggregate bandwidth must scale to thousands of GPUs.

The access-pattern decision becomes concrete when matched to ML workflow stages. For online feature serving, the high-IOPS characteristics of databases enable millisecond lookups of individual records. Large recommendation systems exemplify this challenge at its most extreme: terabyte-scale lookup data may need to serve billions of sparse reads per second, requiring storage architectures that optimize IOPS over sequential throughput. More generally, a recommendation system looking up a user's profile during real-time inference requires random access optimized for per-request latency.

Structured model training points the decision in the opposite direction: throughput dominates request-level latency. The throughput-optimized design of data warehouses enables high-speed sequential scans over large, clean tables. Training a fraud detection model that processes millions of transactions with hundreds of features per transaction benefits from columnar storage that reads only relevant features efficiently, directly reducing the data term by minimizing bytes transferred.

Exploratory analysis and unstructured training data add a third constraint: the schema may not be known when the data is collected. For images, audio, and text, data lakes provide the flexibility and low-cost storage needed for massive volumes. A computer vision system storing terabytes of raw images alongside metadata, annotations, and intermediate processing results benefits from the schema flexibility and cost efficiency that data lakes commonly provide, while the scale of the D_{vol} numerator demands high aggregate bandwidth.

Databases earn their place when transactional consistency and point-lookup latency dominate. They maintain product catalogs, user profiles, or transaction histories with strong consistency guarantees and low-latency point lookups. For ML workflows, databases serve specific roles well: storing feature metadata that changes frequently, managing experiment tracking where transactional consistency matters, or maintaining model registries that require atomic updates. A PostgreSQL database handling structured user attributes (`user_id`, `age`, `country`, `preferences`) provides millisecond lookups for serving systems that need individual user features in real-time. However, databases struggle when ML training requires scanning millions of records repeatedly across multiple epochs. The row-oriented storage that optimizes transactional lookups becomes inefficient when training needs only 20 of 100 columns from each record but must read entire rows to extract those columns.

Data warehouses earn their place when repeated scans over structured features dominate. Columnar storage formats (Stonebraker et al. 2018) enable reading specific features without loading entire records, essential when tables contain hundreds of columns but training needs only a subset. The

I/O reduction scales with the fraction of columns selected: reading 20 of 100 columns transfers about one fifth of the uncompressed row payload before compression or metadata overhead. The format-efficiency calculation in section 4.8.2 quantifies this gain with a worked fraud-detection example. Many successful ML systems draw training data from warehouses because the structured environment simplifies exploratory analysis and iterative development. Data analysts can quickly compute aggregate statistics, identify correlations between features, and validate data quality using familiar SQL interfaces.

However, warehouses assume relatively stable schemas and struggle with truly unstructured data (images, audio, free-form text) or rapidly evolving formats common in experimental ML pipelines. When a computer vision team wants to store raw images alongside extracted features, multiple annotation formats from different labeling vendors, intermediate model predictions, and learned representation vectors, forcing all these into rigid warehouse schemas creates more friction than value. Schema evolution becomes painful: adding new feature types requires ALTER TABLE operations that may take hours on large datasets, blocking other operations and slowing iteration velocity.

Data lakes earn their place when schema flexibility and low-cost retention dominate. They address warehouse limitations by storing structured, semi-structured, and unstructured data in native formats, deferring schema definitions until the point of reading, a pattern called schema-on-read.³²

This flexibility proves valuable during early ML development when teams experiment with diverse data sources and are not yet certain which features will prove useful. A recommendation system might store in the same data lake: transaction logs as JSON, product images as JPEGs, user reviews as text files, clickstream data as Parquet, and model embeddings as NumPy arrays. Rather than forcing these heterogeneous types into a common schema upfront, the data lake preserves them in their native formats. Applications impose schema only when reading, enabling different consumers to interpret the same data differently: one team extracts purchase amounts from transaction logs while another analyzes temporal patterns, each applying schemas suited to their analysis.

That flexibility is only useful if governance prevents the lake from becoming opaque. Without disciplined metadata management and cataloging, data lakes degrade into “data swamps,” disorganized repositories where finding relevant data becomes nearly impossible, undermining the productivity benefits that motivated their adoption. A data lake might contain thousands of datasets across hundreds of directories with names like `userdata_v2_final` and `userdata_v2_final_ACTUALLY_FINAL`, where only the original authors (who have since left the company) understand what distinguishes them. Successful data lake implementations maintain searchable metadata about data lineage, quality metrics, update frequencies, ownership, and access patterns, essentially providing warehouse-like discoverability over lake-scale data. Tools like AWS Glue Data Catalog, Apache Atlas, or Databricks Unity Catalog provide this metadata layer, enabling teams to discover and understand data before investing effort in processing it.

Each storage architecture favors a different access pattern and ML workflow stage, as table 4.7 makes explicit. A poor match can impose substantial performance penalties.

Table 4.7: Storage System Characteristics: Each storage architecture favors a different access pattern and ML workflow stage. Databases support the high-IOPS random access needed for real-time feature serving; data warehouses provide sequential throughput for batch training and feature engineering; and data lakes provide schema flexibility and cost efficiency for large-scale raw data retention. A poor match between system and workload can impose substantial performance penalties.

Attribute	Conventional Database	Data Warehouse	Data Lake
Purpose	Operational and transactional	Analytical and reporting	Storage for raw and diverse data for future processing
Data type	Structured	Structured	Structured, semi-structured, and unstructured
Scale	Small to medium volumes	Medium to large volumes	Large volumes of diverse data
Performance Optimization	Optimized for transactional queries (OLTP)	Optimized for analytical queries (OLAP)	Optimized for scalable storage and retrieval
Examples	MySQL, PostgreSQL, Oracle DB	Google BigQuery, Amazon Redshift, Microsoft Azure Synapse	Google Cloud Storage, AWS S3, Azure Data Lake Storage

³² | **Schema-on-read:** Applies data structure definitions at query time rather than during ingestion, contrasting with schema-on-write (traditional databases) where data must conform to a predefined structure before storage. For ML pipelines in early development, schema-on-read enables rapid experimentation – teams can store raw sensor data, images, and logs without committing to a feature schema upfront. The trade-off is governance: without enforced schemas, data lakes degrade into “data swamps” where finding and validating training data becomes the bottleneck instead.

Choosing appropriate storage requires evaluating workload requirements rather than following technology trends. The decision typically follows a maturity trajectory: early-stage projects start with databases (familiar SQL, existing infrastructure), migrate to warehouses when analytical queries overwhelm transactional performance, and adopt data lakes when unstructured data types (images, audio, text) or petabyte-scale cost optimization become critical. Mature ML organizations typically employ all three, orchestrated through unified data catalogs: databases for operational data and real-time serving, warehouses for curated analytical data and feature engineering, and data lakes for raw heterogeneous data and large-scale training. Consider a self-driving car system: vehicle telemetry lives in a database for real-time monitoring, aggregated driving statistics reside in a warehouse for batch analytics, and terabytes of raw camera images and lidar point clouds occupy a data lake for model training—each storage tier optimized for its access pattern.

4.8.2 Storage performance and cost

Beyond the functional differences between storage systems, cost and performance characteristics directly impact ML system economics and iteration speed. Understanding these quantitative trade-offs enables informed architectural decisions based on workload requirements.

Table 4.8 shows why ML systems use tiered storage. Under these 2024 assumptions, storing our KWS training dataset (748.8 GB) in object storage costs \$17.2/month, enabling affordable raw-audio retention, while working datasets on Non-Volatile Memory Express (NVMe)³³ cost \$74.9/month—\$224.6/month for active training but load 50× faster.

Table 4.8: Illustrative Storage Cost-Performance Trade-Offs (2024): These planning ranges vary by provider, region, configuration, and workload. Training data loading requires high-throughput sequential access; section D.3.3 establishes the memory hierarchy this pattern must align with. Online serving needs low-latency random reads, while archival storage prioritizes cost over access speed for compliance and historical data.

Storage Tier	Cost (\$/TB/month)	Sequential Read Throughput	Random Read Latency	Typical ML Use Case
NVMe SSD (local)	\$100–300	5–7 GB/s	10–100 μs	Training data loading, active feature serving
Object Storage (S3, GCS)	\$20–25	100–500 MB/s (per connection)	10–50 ms	Data lake raw storage, model artifacts
Data Warehouse (BigQuery, Redshift)	\$20–40	1–5 GB/s (columnar scan)	100–500 ms (query startup)	Training data queries, feature engineering
In-Memory Cache (Redis, Memcached)	\$500–1000	20–50 GB/s	1–10 μs	Online feature serving, real-time inference
Archival Storage (Glacier, Nearline)	\$1–4	10–50 MB/s (after retrieval)	Hours (retrieval)	Historical retention, compliance archives

The performance difference directly impacts iteration velocity. Training that loads data at 5 GB/s completes dataset loading in 149.8 s, compared to 7,488 s at typical object storage speeds. This 50× difference determines whether teams can iterate multiple times daily or must wait hours between experiments.


The latency gap between on-chip registers and remote storage spans over eight orders of magnitude, making unbuffered I/O requests catastrophic for GPU accelerator utilization. Scaling these microsecond delays to intuitive human timeframes (table 4.9) makes the penalty of an un-cached fetch immediately visceral.

These latency numbers were popularized by Jeff Dean’s influential 2009 LADIS keynote.³⁴ They explain why a poorly designed storage architecture can leave an expensive accelerator idle, and why distributed training requires careful attention to data locality.

Table 4.9: Memory and Storage Latency Hierarchy: Standard hardware access times normalized to a human scale where a 0.5 ns L1 cache access represents one second.

Operation	Latency (ns)	Human Scale	ML System Impact
L1 Cache Reference	0.5	1 second	Immediate

33 | **NVMe (non-volatile memory express):** A storage protocol connecting directly to the PCIe bus with 64K command queues, delivering 5–7 GB/s sequential throughput and microsecond-scale latency. The contrast with SATA SSD (500 MB/s, single queue) is a 10× bandwidth gap that can bottleneck a training pipeline when storage service time exceeds compute time.



Access latency spans eight orders of magnitude, cache to internet.

34 | **Jeff Dean:** Google Senior Fellow, architect of MapReduce, BigTable, Spanner, and TensorFlow. His 2009 LADIS keynote distilled the numbers in table 4.9 into the engineering heuristic that an L1 cache reference (0.5 ns) and a cross-data center round trip (150 ms) span a 3×10^8 ratio – eight orders of magnitude that explain why a training pipeline reading features from remote storage instead of local NVMe starves the accelerator it was meant to feed.

Operation	Latency (ns)	Human Scale	ML System Impact
L2 Cache Reference	7	14 seconds	Fast computation
Main Memory (DRAM)	100	3 minutes	The “memory wall” threshold
SSD (local NVMe)	100,000	2 days	Data loading bottleneck
Network (same DC)	500,000	11.57 days	Distributed coordination lag
SSD (remote network)	2,000,000	46.30 days	Training-serving skew source
Object Store (S3)	20,000,000	1 year	Archival access
Internet (CA to VA)	100,000,000	6 years	Global user experience

Access pattern alone does not exhaust the storage problem. ML workloads also store artifacts that conventional databases and warehouses were not designed around, and those artifacts shape infrastructure decisions across the entire development lifecycle, from experimental notebooks to production serving systems handling millions of requests per second.

Modern ML models contain millions to trillions of parameters requiring storage and retrieval patterns fundamentally different from traditional data. GPT-3 (Brown et al. 2020) requires approximately 700 GB for model weights when stored in FP32 format (175B parameters times 4 bytes), though practical deployments often use smaller numeric formats such as FP16 (350 GB) to reduce storage and access cost. Even at FP16 precision, this exceeds many organizations’ entire operational databases. The trajectory reveals accelerating scale: from AlexNet’s 60M parameters (Krizhevsky et al. 2012) to GPT-3’s 175B parameters (Brown et al. 2020), model size grew approximately 2916× in eight years. Storage systems must handle these dense numerical arrays efficiently for both capacity and access speed. Unlike typical files where sequential organization matters for readability, model weights benefit from block-aligned storage enabling parallel reads across parameter groups. When multiple accelerators need to read model data from shared storage, whether during training initialization or checkpoint loading, storage systems must deliver aggregate bandwidth approaching network interface limits, often 25 Gbps or higher, without introducing bottlenecks that would idle expensive compute resources. Chapter 10 later explains the model-side techniques that reduce these footprints further.

The iterative nature of ML development introduces versioning requirements qualitatively different from traditional software. Git excels at tracking code changes where files are predominantly text with small incremental modifications, but it fails for large binary files where even small model changes result in entirely new checkpoints. Storing ten versions of a 10 GB model naively would consume 100 GB; ML versioning systems typically keep lightweight metadata pointers in Git or a registry and store artifacts in external content-addressed storage. Identical files can be deduplicated, but changed binary checkpoints are often stored as separate objects unless the backend adds its own delta compression. Tools like DVC (Data Version Control) and MLflow maintain pointers to model artifacts rather than storing copies, enabling efficient versioning while preserving the ability to reproduce any historical model. A typical ML project generates hundreds of model versions during hyperparameter tuning—one version per training run as engineers explore learning rates, batch sizes, architectures, and regularization strategies. Without systematic versioning capturing training configuration, accuracy metrics, and training data version alongside model weights, reproducing results becomes impossible when yesterday’s model performed better than today’s but teams cannot identify which configuration produced it. This reproducibility challenge connects directly to the governance requirements addressed by section 4.8.4: regulatory compliance often requires demonstrating exactly which data and process produced specific model predictions.

Large-scale training generates substantial intermediate data requiring storage systems to handle concurrent read/write operations efficiently. When training jobs use multiple accelerators, each processing unit works on different portions of data, requiring storage systems to handle many simultaneous reads and writes. The specific patterns depend on the parallelization strategy employed, which chapter 8 examines in detail. From a storage perspective, systems must handle concurrent I/O at rates proportional to the number of processing units, with each potentially writing tens to hundreds of megabytes of intermediate results during model updates. For now, treat checkpoint files, optimizer-state snapshots, and explicit offload paths as large binary artifacts produced during training. Storage systems must provide low-latency access to support efficient coordination. If

workers spend more time waiting for storage than performing computations, parallel processing becomes counterproductive regardless of the specific training approach used.

The bandwidth hierarchy that drove the coordination tax in section 4.6.3 constrains ML system design at every level, creating bottlenecks that no amount of compute optimization can overcome. While RAM delivers 50 to 200 gigabytes per second bandwidth on modern servers, network storage systems typically provide only one to ten gigabytes per second, and even high-end NVMe SSDs max out at one to seven gigabytes per second sequential throughput. Modern GPUs can process data faster than storage can supply it, creating scenarios where expensive accelerators idle waiting for data. Consider training an image classification model: loading 1,000 images per second at 150 KB each requires 150 MB/s sustained throughput from storage. When the GPU can process images faster than storage delivers them, the *data pipeline*, not the *model*, becomes the bottleneck. A 10-fold mismatch between GPU processing speed and storage bandwidth means expensive accelerators sit idle 90 percent of the time waiting for data. No amount of GPU optimization can overcome this I/O constraint.

Understanding these quantitative relationships enables informed architectural decisions about storage system selection and data pipeline optimization, which become even more critical during distributed training. Training throughput is bounded by the minimum of compute capacity and data supply rate: when storage cannot keep the accelerator fed, the bottleneck shifts from silicon to I/O.

Napkin Math 4.9: Storage bandwidth budget

Problem: What storage configuration keeps a reference accelerator from being data-starved while training ResNet-50? Chapter 11 explains the device family later; here, the point is how a compute ceiling becomes a data-path budget.

Start with the compute ceiling. The reference accelerator can deliver 312 TFLOP/s on dense FP16 operations. ResNet-50 costs about eight GFLOPs for each forward pass, and including the backward training pass raises the training-step cost to about 24.6 GFLOP per image. The training chapter derives that backward-pass machinery; here, the combined per-image cost is the input to the storage budget. When we divide accelerator peak by model cost, the calculation gives us an upper bound of 12,682 img/s.

That compute ceiling becomes a storage requirement once each image must arrive from disk or object storage. With 150 KB JPEG-compressed images, the data path must supply that many images per second times 150 KB per image, or approximately 1.9 GB/s.

The storage options now have a concrete target: saturating this accelerator requires 1.9 GB/s of sustained bandwidth.

- S3 Standard delivers about 100 MB/s per thread, so the pipeline needs 19 concurrent worker threads before software overhead.
- SATA SSDs deliver about 500 MB/s sequentially, making them a bottleneck for this accelerator.
- NVMe SSDs deliver approximately 5–7 GB/s, which is the right class of local storage for the target.

Systems insight: With SATA SSDs, maximum throughput is capped by 500 MB/s divided across 150 KB images, or approximately 3,333 img/s. The \$15,000 GPU will run at 26 percent utilization because storage supplies only a small fraction of the accelerator’s image-processing ceiling. Storage physics dictates training speed.

Designing for high-throughput training starts by matching storage throughput to accelerator demand. This is the chapter’s starvation argument rotated to its third and final lens: section 4.2.3 priced the idle accelerator as a feeding tax, section 4.5.5 traced the same stall to CPU decode workers, and here the question becomes which storage tier can sustain the required supply rate.

The 500 MB/s figure represents effective SATA III sequential read throughput (the interface ceiling is 550 MB/s), and real-world random read performance with small files can be significantly

lower. That caveat reinforces the general principle governing data pipelines: equation 4.5 bounds training throughput by compute capacity and the data supply rate defined in equation 4.6. Let R_{train} , R_{compute} , and R_{data} denote rates in samples per second; let B_{storage} be storage bandwidth in bytes per second, η_{overhead} the dimensionless fraction lost to overhead, and S_{sample} the bytes per sample. Under full overlap, the min-of-rates form corresponds to the $T_{\text{step}} = \max(T_{\text{compute}}, T_{\text{io}})$ lower bound from section 4.5.5; with incomplete overlap, actual step time is higher.

$$R_{\text{train}} \leq \min(R_{\text{compute}}, R_{\text{data}}) \quad (4.5)$$

$$R_{\text{data}} = \frac{B_{\text{storage}}(1 - \eta_{\text{overhead}})}{S_{\text{sample}}} \quad (4.6)$$

Once R_{data} is the smaller rate, additional accelerator compute cannot raise training throughput; the remedy must increase usable storage bandwidth or reduce bytes per sample.

When storage bandwidth becomes the limiting factor, teams must either improve storage performance through faster media, parallelization, or caching, or reduce the amount of data that must move. Large language model training may require processing hundreds of gigabytes of text per hour, while computer vision models processing high-resolution imagery can demand sustained data rates exceeding 50 gigabytes per second across distributed clusters. These requirements make data loading a systems-placement decision: framework data loaders parallelize I/O across asynchronous worker processes, using **double-buffering prefetch queues** to overlap the $D_{\text{vol}}/BW_{\text{IO}}$ data fetch phase of batch $k+1$ in host RAM while the accelerator computes the O/R_{peak} pass of batch k on GPU memory, preventing compute starvation stalls (L_{lat}). Framework loaders can also move expensive augmentation work closer to the accelerator rather than storing every augmented variant.

File format selection dramatically impacts the data term ($\frac{D_{\text{vol}}}{BW}$) of the iron law. Columnar formats like **Parquet** excel for tabular features by allowing column projection and pushdown filtering, whereas shard-based sequential formats like **WebDataset** or **TFRecord** pack millions of small binary objects (e.g., images or audio) into contiguous sequential streams. This sharding eliminates random disk seek overhead (L_{lat}), transforming slow random I/O into sustained sequential storage throughput (BW_{IO}). We can quantify this impact as *format efficiency* (η_{format}), which acts as a multiplier on effective bandwidth.

Napkin Math 4.10: Format efficiency

Storage formats determine how much “waste” data must move to retrieve the signal needed for training. Let η_{format} be the dimensionless fraction of bytes read that carry selected features, from 0 to 1. Equation 4.7 shows how that fraction scales physical bandwidth into usable pipeline bandwidth:

$$\text{Effective Bandwidth} = \text{Physical Bandwidth} \times \eta_{\text{format}} \quad (4.7)$$

Values below 1 mean the format spends part of the physical data path on irrelevant bytes; improving layout can therefore raise usable throughput without changing the storage device.

Scenario: Training a fraud model using 20 features from a 100-column table.

Row-oriented (CSV) storage must read all 100 columns to get the 20 needed, giving a useful-byte fraction of 0.2 and wasting 80 percent of disk bandwidth. **Column-oriented (Parquet)** storage reads only the needed columns, so $\eta_{\text{format}} \approx 1$ ignoring metadata overhead, and the pipeline gets 5× higher effective throughput.

Systems insight: In this projection-only scenario, switching from CSV to Parquet provides the same effective read-throughput gain as a 5× faster data path. Section B.1.3 treats row vs. columnar storage layouts and the algebra of data operations (selection, projection, join) in depth.

Format choice follows the workload’s access pattern. When a large tabular dataset cannot be moved cheaply and training reads only selected columns, a columnar format reduces the bytes read per step.

Columnar storage formats like Parquet or Optimized Row Columnar (ORC) realize this reduction, five to ten times for typical ML workloads, through two mechanisms: the column projection the format-efficiency calculation just quantified, and column-level compression exploiting value patterns within columns. Column compression proves particularly effective for categorical features with limited cardinality: a country code column with 200 unique values in 100 million records compresses 20 to 50 times through dictionary encoding, while run-length encoding compresses sorted columns by storing only value changes. The combination can achieve total I/O reduction of 20 to 100 times compared to uncompressed row formats, directly translating to faster training iterations and reduced infrastructure costs.

Compression algorithm selection involves trade-offs between compression ratio and decompression speed. While gzip achieves higher compression ratios of six to eight times, Snappy achieves only two to three times compression but decompresses at 500 MB/s, roughly 4.2× faster than gzip's 120 MB/s. For ML training where throughput matters more than storage costs, Snappy's speed advantage often outweighs gzip's space savings. Training on a 100 GB dataset compressed with gzip requires 13.9 minutes of decompression time, while Snappy requires only 3.3 minutes. When training iterates over data for 50 epochs, this 10.6 minutes difference per epoch compounds to 9 hours total, potentially the difference between running experiments overnight vs. waiting multiple days for results. The choice cascades through the system: faster decompression enables higher input throughput, reduced buffering requirements (less decompressed data needs staging), and better GPU utilization (less time idle waiting for data).

Storage performance optimization extends beyond format and compression to data layout strategies. Data partitioning based on frequently used query parameters dramatically improves retrieval efficiency. A recommendation system processing user interactions might partition data by date and user demographic attributes, enabling training on recent data subsets or specific user segments without scanning the entire dataset. Partitioning strategies interact with distributed training patterns: range partitioning by user ID enables data parallel training where each worker processes a consistent user subset, while random partitioning ensures workers see diverse data distributions. The partitioning granularity matters: too few partitions limit parallelism, while too many partitions increase metadata overhead and reduce efficiency of sequential reads within partitions. Poor partitioning compounds these problems in multi-accelerator training: when a distributed data-parallel job assigns shards to workers, imbalanced partition sizes cause straggler effects where the slowest reader holds up the entire gradient synchronization step. If all workers in an eight-accelerator job finish loading their batch in 12 ms but one slow worker reads from an oversized partition and takes 180 ms, that straggler determines the effective batch time. Well-partitioned datasets where each shard fits a worker's local read budget and can be fetched without contending on a shared file handle are therefore a prerequisite for keeping distributed accelerator utilization high, a concern examined further in chapter 8.

4.8.3 Storage across the ML lifecycle

Storage requirements evolve because each lifecycle stage asks the same data to serve a different access pattern. The same dataset is accessed through random sampling during exploratory analysis, sequential scanning during model training, and random access during production serving. These diverse patterns require storage architectures that accommodate all three access modes.

During development, flexibility matters more than raw performance. The key challenge is managing dataset versions without overwhelming storage capacity: ten experiments on a 100 GB dataset would naively require 1 TB of copies. The metadata-pointer mechanism introduced for model checkpoints in section 4.8.2 applies unchanged: tools like DVC track versions through pointers and content-addressed artifact storage, deduplicating identical content when possible. Governance considerations demand tiered access controls where synthetic or anonymized datasets are broadly available for experimentation, while production data containing sensitive information requires approval and audit trails.

Training phase requirements shift dramatically toward throughput. Modern deep learning processes massive datasets repeatedly across dozens or hundreds of epochs, making I/O efficiency critical. Training ResNet-50 on ImageNet across eight GPUs at 40,000 img/s would require roughly 6 GB/s for 150 KB compressed images, and substantially more if decoded FP32 tensors are staged.

Storage unable to sustain this throughput idles GPUs, directly increasing infrastructure costs. The feature computation placement trade-off (section 4.5.5.3) is especially acute here: precomputing features achieves 30× storage reduction (150 KB to 5 KB vectors) but introduces staleness risk when extraction logic changes.

Deployment and real-time inference requirements prioritize low-latency random access. A recommendation system serving 10,000 req/s with 10 ms latency budgets and 10 feature reads/request requires 100,000 IOPS; meeting that illustrative target may require an in-memory database such as Redis, aggressive caching, or a purpose-built distributed key-value service. Edge deployment adds further constraints: limited device storage, intermittent connectivity, and the need for model updates without disrupting inference, typically addressed through tiered storage where models cache locally while reference data pulls from the cloud. Model versioning must support smooth transitions between versions, rapid rollback, and serving multiple versions simultaneously for A/B testing, operational patterns examined in chapter 14. Those serving and rollback guarantees depend on provenance: the system must know the exact dataset version behind each model version.

4.8.4 Data versioning for ML reproducibility

Suppose a recommendation model’s click-through rate drops after a routine weekly retraining even though the code has not changed. The team must distinguish among a data change, a labeling shift, and a corrupted upstream table. Without a record linking the model to the exact dataset that produced it, the team may spend days bisecting possibilities. With data versioning, they can diff the training snapshots and identify an upstream backfill that shifted the label distribution. Listing 4.3 shows the mechanism: Git records the small pointer file, DVC moves the large data bytes to remote storage, and a later checkout restores the exact training snapshot paired with the code commit (Iterative 2024).

Listing 4.3: DVC Workflow: Git-like semantics for data versioning. DVC tracks large data files alongside code commits, enabling exact reproduction of any historical training dataset through paired `git checkout` and `dvc checkout` commands.

```
# Add data to version control
dvc add data/training.csv
git add data/training.csv.dvc
git commit -m "Add training data v1"
dvc push # Upload to remote storage

# Later: retrieve exact data for any historical commit
git checkout abc123
dvc checkout # Restores exact data from that commit
```

Data versioning is the storage analogue of source-control provenance. It connects model versions to exact training data, enabling debugging and reproducibility. Without it, teams cannot identify the exact data that trained the model now misbehaving in production.

DVC (Data Version Control) provides Git-like semantics for file snapshots (Iterative 2024), while listing 4.4 demonstrates Delta Lake’s transaction-log-based historical queries (Armbrust et al. 2020).

Listing 4.4: Delta Lake Time Travel: SQL queries identify a historical table state by timestamp or transaction-log version.

```
-- Query data as it existed on a specific date
SELECT * FROM training_data TIMESTAMP AS OF '2024-01-15'

-- Or by version number for programmatic access
SELECT * FROM training_data VERSION AS OF 47
```

Two complementary capabilities complete the versioning infrastructure. Feature store point-in-time retrieval maintains historical feature values, enabling training with features “as they existed” at prediction time and preventing label leakage, the accidental use of information that would not be available when the prediction is made. **Model registry** integration links each model registry entry, a

catalog record for a model artifact and its metadata, to its provenance: Git commit hash, data version, feature snapshot timestamp, and training configuration. Applied to the click-through-rate drop that opened this section, complete records can turn an open-ended search into a direct snapshot comparison that exposes the upstream backfill.

Long-term maintenance introduces a final storage consideration: retaining enough data to debug issues while honoring privacy, cost, and applicable record-keeping requirements. A high-volume recommendation system may use tiered retention: hot storage for rapid analysis, warm storage for periodic review, and cold archive storage for approved long-term investigations or records. Retention periods and immutability requirements depend on the data, jurisdiction, product, and risk; keeping everything indefinitely is neither a universal compliance rule nor a sound default.

Storage architecture determines where data resides and how it is retrieved, but physical storage layout alone cannot guarantee that features computed during batch training match those served in real time. Bridging batch training environments and low-latency serving systems requires specialized infrastructure designed for dual-access consistency.

4.8.5 Feature stores

The fundamental challenge of production feature management is preserving semantic consistency across environments: serving historical feature values for training and current feature values for inference. **Point-in-time correctness** is one reason teams adopt **feature stores**, which can reduce training-serving skew and enable feature reuse across models and teams. They do not eliminate skew by themselves; timestamp handling, backfills, freshness, and offline-online synchronization must still be correct.

The core problem feature stores address becomes clear when examining typical ML development workflows. During model development, data scientists write feature engineering logic in notebooks or scripts, often using different libraries and languages than production serving systems. Training might compute a user's "total purchases last 30 days" using SQL aggregating historical data, while serving computes the same feature using a microservice that incrementally updates cached values. These implementations should produce identical results, but subtle differences in handling timezone conversions, dealing with missing data, or rounding numerical values cause training and serving features to diverge. Uber's Michelangelo platform description treats training-serving skew and reusable feature pipelines as central production concerns, motivating an integrated platform approach to feature management (Hermann and Del Balso 2017).



Definition 4.4: Feature store

Feature store is the architectural layer that centralizes the management of machine learning features, decoupling feature computation from consumption.

1. **Significance:** It supports point-in-time-correct historical retrieval for training ($x_{t-\Delta}$) and coordinated feature definitions for real-time inference (x_t), reducing training-serving skew when timestamps, backfills, freshness, and synchronization are handled correctly.
2. **Distinction:** Unlike a general-purpose database, a feature store is designed for dual storage modes: an *offline store* (columnar/batch) for training and an *online store* (key-value/low-latency) for serving.
3. **Common pitfall:** A frequent misconception is that a feature store is only "a place to store data." In reality, it manages feature definitions, values, and retrieval; transformation may run in the feature store, an offline store, or a separate batch or stream-compute engine.

Feature stores (discussed architecturally in section 14.4.1.2) provide a shared source of truth for feature definitions, supporting consistency across the ML lifecycle. When data scientists define a feature like `user_purchase_count_30d`, the feature store can maintain the definition (SQL query, transformation logic, or computation graph) and provide historical values for training and current values for serving. This architectural pattern reduces a class of subtle bugs that prove notoriously difficult to debug because models train successfully but perform poorly in production without obvious errors. The same centralized approach enables feature reuse across models and teams:

when multiple teams build models requiring similar features, the feature store can prevent each team from reimplementing identical computations with subtle variations. A recommendation system might compute user embedding vectors across hundreds of dimensions, aggregating months of interaction history. Rather than each model team recomputing these expensive embeddings, the feature store can compute them once and serve them to multiple consumers.

The architectural pattern typically implements dual storage modes optimized for different access patterns. The offline store uses columnar formats like Parquet on object storage, optimized for batch access during training where sequential scanning of millions of examples is common. The online store uses key-value systems like Redis, optimized for random access during serving where individual feature vectors must be retrieved in milliseconds. Synchronization between stores becomes critical. As training generates new models using current feature values, those models deploy to production expecting the online store to serve consistent features. Feature stores typically implement scheduled batch updates propagating new feature values from offline to online stores, with update frequencies depending on feature freshness requirements.

Time-travel capabilities distinguish sophisticated feature stores from simple caching layers. Training requires accessing feature values as they existed at specific points in time rather than current values. Consider training a churn prediction model: for users who churned on January 15th, the model should use features computed on January 14th, not current features reflecting their churned status. Point-in-time correctness ensures training data matches production conditions where predictions use currently-available features to forecast future outcomes. Implementing time-travel requires storing feature history, not just current values, substantially increasing storage requirements but enabling correct training on historical data.

Feature store performance characteristics directly impact both training throughput and serving latency. The offline store must support high-throughput batch reads (millions of feature vectors per minute) using columnar formats that enable efficient reads of specific features from wide tables. The online store must support thousands to millions of reads per second with single-digit millisecond latency. In production, feature freshness adds further pressure: when users add items to shopping carts, recommendation systems need updated features within seconds, not hours. Streaming feature computation pipelines address this by updating online stores continuously rather than through periodic batch jobs, though streaming introduces complexity around exactly-once processing semantics—ensuring each event updates state once despite retries—and handling late-arriving events.

A fully assembled pipeline covering acquisition, ingestion, processing, labeling, and storage might suggest that data engineering work is “done.” Production systems, however, do not stand still. User behavior drifts, upstream schemas evolve, labeling guidelines change, and the pipeline engineering described in earlier sections gradually erodes unless actively maintained.

4.9 Operational Data Health

Imagine a KWS system that launched with 98 percent accuracy and excellent training-serving consistency. Six months later, accuracy has slipped to 94 percent, not because anyone changed the model, but because new smartphone microphone hardware shifted the audio distribution, labeling guidelines were updated without retraining, and a schema change in the user metadata pipeline went undocumented. Each is a form of **data debt**: accumulated compromises in data quality, documentation, and infrastructure that compound silently. Data debt can affect model performance and remain invisible until failures become severe. Managing that debt requires classifying it, measuring its accumulation, and debugging the pipeline when the debt comes due.

4.9.1 Categories of data debt

The KWS scenario in the preceding section illustrates each of the four categories of data debt. The microphone hardware shift is *freshness debt*; the updated labeling guidelines are *quality debt*; the undocumented schema change is both *schema debt* and *documentation debt*. Each category requires different detection and remediation strategies, but all share a common structure.

**Definition 4.5: Data debt**

Data debt is the compound interest of implicit coupling and missing documentation across an ML data stack.

1. **Significance:** It can manifest as silent degradation and rising maintenance costs caused by unmanaged dependencies or unaddressed distribution shifts.
2. **Distinction:** Unlike code-focused technical debt, data debt centers on data assumptions, quality, provenance, and distribution; both can affect development cost, reliability, and model behavior.
3. **Common pitfall:** A frequent misconception is that data debt is just “bad data.” In reality, it is a systems architecture failure: it occurs when the assumptions of the training distribution are no longer enforced at the system boundary.

The first category is **documentation debt**: data provenance, meaning, and quality characteristics go unrecorded. Datasets without datasheets or data cards (Gebru et al. 2021; Pushkarna et al. 2022), unlabeled columns, and undocumented transformations create debt that compounds when original authors leave organizations. Sambasivan et al. (2021) describe missing metadata, undocumented assumptions, and poor cross-organizational documentation as contributors to data cascades in high-stakes AI systems. The symptoms are unmistakable: columns whose meanings must be reverse-engineered from usage patterns, missing provenance records that block compliance audits, undocumented assumptions that cause silent failures when those assumptions change, and absent quality metrics that prevent informed dataset selection.

The second category, **schema debt**, emerges from accumulated schema workarounds and migrations. When upstream systems change data formats, quick fixes such as string parsing instead of proper type handling or NULL coercion instead of error handling accumulate into fragile transformation logic. Telltale patterns include multiple date handlers for one logical field, defensive null checks scattered throughout pipeline code, version-specific parsing branches, and undocumented enum mappings. Depending on the pipeline, an unexpected column may be ignored, fail validation, or change the assembled tensor shape; an unseen categorical value may map to an unknown token or exceed an embedding vocabulary. The pipeline jungle example in section 4.3.2 illustrates the general lesson: contract violations can surface far downstream and resemble model bugs.

The third category, **quality debt**, consists of known data errors that remain uncorrected due to time or resource constraints. A dataset with 3 percent known label errors represents quality debt: each training run on this data produces models carrying those errors. Quality debt compounds through a particularly dangerous feedback loop, because models trained on erroneous data make predictions that become training data for downstream systems, amplifying the original errors across the ecosystem. Concrete indicators include known label error rates not yet corrected, documented biases not yet mitigated, identified duplicates not yet deduplicated, and detected drift not yet addressed through retraining.

The fourth category, **freshness debt**, arises when training data diverges from production distributions over time. A model trained on 2023 user behavior deployed in 2025 carries freshness debt: the distribution shift between training and serving can change performance. Freshness debt is particularly dangerous because it accumulates silently. Unlike code that breaks visibly, stale models can degrade gradually until performance drops below acceptable thresholds. Warning signs include time since last retraining exceeding the expected refresh cadence, feature staleness from cached features not updated at the required frequency, and reference data lag where lookup tables no longer reflect current state. Classifying debt is useful only when each category has measurable warning thresholds.

4.9.2 Measuring and projecting data debt

Unlike technical debt, which can be assessed through code complexity metrics, data debt requires specialized measurement approaches. Table 4.10 gives example warning and critical thresholds for each debt category, making data debt trackable as an engineering service-level objective (SLO) rather than presenting universal standards.

Table 4.10: Data Debt Metrics: Illustrative thresholds for detecting data debt accumulation across four categories. Warning thresholds indicate debt requiring attention in upcoming planning cycles; critical thresholds indicate debt requiring immediate remediation to prevent system degradation. The PSI row uses common operational drift-monitoring bands, while the remaining rows should be calibrated to the risk tolerance and failure cost of the specific ML system.

Debt Category	Metric	Warning Threshold	Critical Threshold
Documentation	% datasets with data cards	< 80%	< 50%
Documentation	% columns with descriptions	< 90%	< 70%
Schema	Schema version branches	> 3	> 10
Schema	% transformations with try/catch	> 20%	> 50%
Quality	Known label error rate	> 1%	> 5%
Quality	Documented bias metrics	Not measured	Measured but not mitigated
Freshness	Days since last retraining	> 90 days	> 365 days
Freshness	PSI vs. training baseline	> 0.1	> 0.25

Data debt can compound through several feedback loops in ML systems. Training amplification can propagate errors when predictions become features or pseudo-labels for downstream systems. Documentation decays as original authors leave and systems evolve, leaving datasets increasingly opaque. Schema workarounds create brittle code, and additional workarounds can increase the risk that future schema changes cause failures. Distribution drift can widen the gap between training and serving distributions when monitoring and model updates lapse, although its effect on model quality depends on the system. To illustrate sensitivity to a sustained accumulation rate, let Debt_0 be the initial debt level, r an assumed rate per period, and n the number of periods. The scenario in equation 4.8 uses:

$$\text{Debt}_n \approx \text{Debt}_0 \times (1 + r)^n \quad (4.8)$$

This equation is a planning model, not an empirical law; actual debt need not follow exponential growth or a constant rate.

4.9.3 Remediation strategies

Tracking data-debt indicators turns a vague maintenance problem into an engineering budget: remediation can be scheduled before the burden pushes the system into crisis repair. Addressing data debt therefore requires systematic investment, not heroic one-time efforts. Each debt category calls for a distinct remediation approach, but all share a common pattern: regular, budgeted effort rather than crisis-driven scrambles.


One example policy for documentation debt is a dedicated sprint each quarter for recording data provenance, column semantics, and transformation logic. Teams can prioritize heavily used datasets first, with the cadence and scope determined by their backlog and risk.

Schema debt requires preventive infrastructure rather than reactive fixes. Explicit schema contracts between data producers and consumers, codified through tools like Great Expectations or Pandera, fail fast when schemas drift rather than allowing workarounds to accumulate. The investment in contract enforcement pays dividends by preventing the brittle parsing logic that characterizes mature systems with unmanaged schema evolution.

Quality debt demands allocated capacity for remediating known label errors, documented biases, and identified duplicates. The appropriate allocation depends on risk and backlog size. The known error backlog should be tracked and its burn-down rate measured like any engineering metric, making quality remediation visible alongside feature development rather than perpetually deferred.

Freshness debt cannot be addressed through periodic heroics; it requires drift alerts and an evidence-based retraining policy, connecting directly to the PSI and KL divergence monitoring infrastructure established in section 4.5.3. The remediation budget must therefore include monitoring coverage and retraining capacity alongside cleanup work after accuracy has already decayed.

Data debt, like technical debt, is not inherently bad. *Strategic* debt (knowingly accepting documentation shortcuts to meet a deadline) can be rational. The danger lies in *unconscious* debt that accumulates untracked until remediation costs exceed system value. Managing this trade-off requires treating data debt not as a moral failing, but as a systems parameter to be optimized alongside latency and throughput.



4.9.4 Debugging data pipelines

When data debt comes due, it surfaces as model accuracy degradation, pipeline failures, or subgroup performance disparities. Effective debugging applies the diagnostic principles established throughout this chapter: data cascades (section 4.3.1) remind us that root causes lie upstream of symptoms, training-serving skew (section 4.6.1) explains many deployment failures, and drift detection (section 4.5.3) surfaces gradual degradation. Figure 4.15 synthesizes these concepts into an actionable diagnostic sequence.




Figure 4.15: Data Pipeline Debugging Flowchart: Four sequential checks prioritize hypotheses involving drift, overfitting or data artifacts, training-serving skew, and subgroup coverage. Each branch suggests a hypothesis to investigate rather than a conclusive diagnosis; if no listed condition holds, the final box suggests examining model architecture or capacity.

Production ML failures frequently trace to data and operational context rather than model architecture alone.³⁵ Training-serving skew, data drift, and label-quality issues are useful early hypotheses, but the flowchart does not exclude model, code, infrastructure, or interaction effects.

4.10 Fallacies and Pitfalls

From acquisition through storage, every pipeline stage introduces opportunities for both excellence and failure. The following fallacies and pitfalls distill the most consequential misconceptions that lead teams astray.

Fallacy: *More data always improves model performance.*

Beyond a threshold, additional data yields diminishing returns. Empirical studies across image classification, translation, and language modeling confirm that test loss often follows a power law in dataset size, so each additional tranche of data produces progressively smaller gains (Hestness et al. 2017). The task-relevant signal heuristic from section 4.2 reflects the same systems trade-off: redundant examples can add data-processing cost while contributing little new task-relevant information.

³⁵ | **ML failure distribution:** Treat the data-dominance pattern as a diagnostic heuristic rather than a universal empirical distribution. Data and operational checks are useful early hypotheses, but they do not establish the cause of a particular failure.

Pitfall: *Planning petabyte migration as a bandwidth-only transfer.*

Engineers may estimate migration time by dividing dataset size by network bandwidth, but data gravity (section 4.2) also pulls in pipeline re-engineering, data-quality revalidation, schema migration, lineage updates, and synchronization of dependent services. Petabyte-scale migrations can therefore require substantially more engineering time than wire time.

Fallacy: *Data preprocessing can be finished once and left alone.*

Data distributions can drift as user behavior, market conditions, and upstream systems evolve. A preprocessing pipeline validated at launch can become stale as the world changes around it. Production systems require monitoring and evidence-based responses, not one-time validation.

Pitfall: *Ignoring data serialization cost.*

Teams can meticulously optimize accelerator kernels while leaving data loading as JSON or CSV. Text formats can add parsing and byte-transfer overhead relative to binary formats, while columnar formats such as Parquet reduce I/O when tabular workloads read only a subset of columns. Benchmark formats against the actual access pattern and convert when the measured savings justify the migration.

Fallacy: *High training accuracy indicates production readiness.*

Training accuracy measures fit to historical data; held-out evaluation estimates generalization under its sampled conditions, while production performance depends on the live environment. Training-serving skew, distribution shift, and coverage gaps can make a model with strong validation results perform poorly after deployment. The debugging flowchart in figure 4.15 exists because diagnosing that gap requires tracing data and system conditions, not trusting one aggregate score.

Pitfall: *Ignoring training-serving skew until deployment.*

Feature-computation differences between training and serving environments are one important cause of ML deployment failures. By the time skew manifests in production metrics, debugging can become difficult. Feature stores and consistency contracts should be considered early rather than only after deployment incidents.

Fallacy: *Synthetic data can fully replace real-world data collection.*

Synthetic data can augment real observations by targeting rare cases, controlled variations, or privacy-sensitive scenarios, but its value depends on fidelity to the deployment task. Generation inherits omissions and artifacts from its models and assumptions. A KWS system trained only on synthesized speech may miss accents, background noises, and pronunciation variations absent from the generator. Teams should validate the real-synthetic mixture against representative deployment data rather than assume a universally optimal recipe.

Pitfall: *Neglecting data versioning until model debugging requires it.*

Teams may defer data versioning until a deployed model produces unexpected results. Without versioning, reproducing a training run requires reconstructing the exact inputs or re-executing the pipeline. Consider a model that performed well three months ago but degrades after retraining on updated data. Without versioned snapshots, the team cannot determine whether the regression stems from a labeling policy change, a schema migration error, or genuine distribution shift. The data lineage principles established in section 4.6.4 formalize this requirement: every training artifact must trace back to a specific, immutable dataset version. Deferring versioning lengthens debugging because each investigation must first identify the exact data that trained the model.

4.11 Summary

Data engineering provides the foundational infrastructure that transforms raw information into the basis of machine learning systems, determining model performance, system reliability, ethical compliance, and long-term maintainability. The four pillars framework of Quality, Reliability, Scalability, and Governance organizes design choices across acquisition, ingestion, validation, and storage, while the cascading nature of data quality failures reveals why every pipeline stage requires careful engineering decisions. The task of “getting data ready” encompasses complex trade-offs

quantified throughout this chapter: data engineering cost constants for budgeting, storage performance hierarchies, and drift detection thresholds that operationalize the degradation equation into production monitoring infrastructure.



Key Takeaways: Data is the source code

- **Data cascades make upstream quality the highest-leverage investment:** Collection errors amplify through every pipeline stage (figure 4.1). The four pillars (Quality, Reliability, Scalability, Governance) organize prevention, while documentation, schema, quality, and freshness debt require continuous remediation.
- **Data is code; version it, test it, review it:** A dataset is source code for an ML system. Apply the same rigor through version control, validation tests, and data review.
- **Training-serving consistency is nonnegotiable:** Feature transformations shared by training and serving must have equivalent semantics and reuse the same learned state, such as normalization constants and vocabulary mappings.
- **Pipeline architecture choices have large cost implications:** Streaming costs more to operate than batch, while ETL trades storage savings for greater schema-change overhead. Select ingestion patterns by the value of latency, not the appeal of real time.
- **Labeling costs dominate and require substantial resource allocation:** Labeling can cost hundreds to more than a thousand times one optimized training run; in the reference calculation in section 4.4.2, the ratio is $521\times$ – $1,562\times$. Labeling can remain a costly scheduling bottleneck even when annotation work is parallelized.
- **Storage hierarchy determines iteration speed:** The $50\times$ throughput gap between local NVMe (5 GB/s) and cloud object storage (100 MB/s) determines whether iterations occur daily or weekly.
- **The degradation equation becomes actionable through drift and outcome monitoring:** Divergence metrics such as KL measure changes between P_t and P_0 , while labeled outcomes or validated proxies determine whether those changes correspond to model degradation.

The KWS case study combines curated, crowdsourced, and synthetic acquisition; consistency validation; tiered storage for 23.4 million audio samples across 748.8 GB of raw data; and lineage for always-listening devices. Data engineering is not preprocessing before “real” ML work; it underpins model performance, user trust, and regulatory compliance.



What’s Next: From source code to executable

The data pipeline has produced a clean, versioned training set, but a dataset teaches nothing on its own. Chapter 5 opens that black box to examine the algorithm axis: the mathematical operations that convert prepared examples into learned behavior and make neural networks predictable, debuggable, and optimizable.

II

DEVELOPMENT AND TRAINING

Development Principles

Building a machine learning system requires more than assembling model components; it requires managing the flow of information and energy through silicon. Part I established that data is both the program and the physical anchor of every ML system. With that anchor in place, Part II turns to the algorithm-machine interaction: *how* mathematical models are co-designed around the physical limits of the hardware that must execute them. The principles here begin with the accounting that determines why certain architectures succeed while others fail at scale.

III Principle 3: The iron law of ML systems

Invariant: Let T denote idealized wall-clock time in seconds. Here D_{vol} is data volume in bytes, BW is the effective bandwidth of the relevant memory or network path in bytes/s, O is total floating-point operations, R_{peak} is peak compute rate in FLOP/s, η_{hw} is dimensionless hardware utilization efficiency, and L_{lat} is fixed latency overhead in seconds, such as a kernel launch or network round trip. When data movement, compute, and fixed overhead execute serially, their idealized times add:

$$T \approx \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}}$$

The full treatment appears in section 1.7.

When data movement and compute overlap on modern hardware, wall-clock time follows the critical path and cannot be shorter than the slower of those stages plus any fixed latency that remains outside the overlap. The practical lesson is about **dominance**, not unconditional summation.

Implication: Optimization is rarely free of trade-offs. Reducing one term often shifts the bottleneck to another. For example, exploiting unstructured sparsity can reduce executed arithmetic, but sparse representations and kernels may lower effective bandwidth or add index traffic, increasing data-movement time (D_{vol}/BW). An optimization improves wall-clock time only to the extent that it shortens the critical path on the target hardware.

The iron law identifies *what* to optimize, but not *how*. The architecture decides *how* before implementation begins.

III Principle 4: The silicon contract

Invariant: Every model architecture and workload regime makes an implicit commitment to the hardware, a wager on which resource it will saturate first.


- **ResNet-50** assumes high-density floating-point compute. In batched training or inference on accelerators, it is often compute bound: performance is limited by $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$.
- **8-billion-parameter Llama 3** assumes high-bandwidth memory access during autoregressive decoding. At small batch sizes, it is often bandwidth bound: performance is limited by D_{vol}/BW .

- **DLRM** assumes massive embedding tables and sparse lookups. It is shaped by both capacity and bandwidth: performance depends on whether embedding tables fit in the available memory hierarchy and how quickly sparse accesses can be served.

Implication: Designing a model without knowing which hardware resource it will saturate is like designing a bridge without knowing the strength of the steel. The design must target the bottleneck.

Together, the iron law and the silicon contract frame every design decision in Part II. The chapters that follow translate these principles into the components of the ML stack: the mathematical foundations of gradient flow, the architectural patterns that commit to specific hardware resources, the frameworks that map model operations onto hardware, and the training systems that execute the same physics at scale.

Neural Computation




- 5.1 From Logic to Arithmetic
- 5.2 Computing with Patterns
- 5.3 Neural Network Fundamentals
- 5.4 Learning Process
- 5.5 Inference Pipeline
- 5.6 Handwritten Digit Recognition
- 5.7 D·A·M Taxonomy
- 5.8 Fallacies and Pitfalls
- 5.9 Summary

Purpose

Why does understanding a neural network's math matter more than reading its code?

Neural network code can look deceptively simple because a few tensor operations hide the workload that every layer of the system must execute. Matrix dimensions determine arithmetic and data movement; intermediate values determine memory pressure; numerical ranges and gradients determine whether learning remains stable. An operation that looks cheap in isolation can dominate a system when repeated across a batch or when its outputs must be retained for learning, so source-code length gives little indication of resource cost. Two implementations of the same model can run very differently because batch shape, precision, and the target machine change what fits, how quickly it runs, and whether its calculations remain numerically valid. When training stalls or deployment exhausts memory, reading a framework call rarely identifies the binding constraint. A numerical failure, a memory bottleneck, and an inefficient implementation can produce similar high-level symptoms while demanding different remedies. Diagnosis and planning therefore require reasoning from the underlying operations to the bytes moved, values retained, and work repeated. This view can expose a mismatch before an expensive training run or deployment attempt does. Those primitives provide a stable vocabulary across architectures and make software abstractions auditable, turning opaque code into a workload that can be measured, diagnosed, and mapped to hardware. In D·A·M terms, neural computation is the interface where algorithmic structure becomes machine work and where physical limits feed back into model design.





Learning Objectives

- Explain why learned arithmetic replaced rule-based logic for high-dimensional pattern recognition
- Calculate parameter counts, multiply-accumulate operations, activation storage, and memory traffic for multilayer networks
- Compare activation functions by nonlinearity, gradient behavior, and hardware execution cost
- Apply cross-entropy loss and gradient updates to reason about supervised learning dynamics
- Derive forward and backward propagation as matrix operations over batches
- Analyze training vs. inference workloads using compute, memory, and deployment constraints
- Evaluate an end-to-end document-recognition pipeline from preprocessing through decision thresholds and human review

5.1 From Logic to Arithmetic

A dataset that has passed the preceding engineering stages is ready for model development, but examples do not learn on their own. The next system boundary is the model, where data becomes learned behavior through matrix multiplications, activation functions, and gradient updates. A model that runs correctly on one machine and fails on another is not necessarily suffering from a hardware defect. Its layer dimensions may create intermediate activations, the per-layer outputs that training retains for the backward pass, that exceed the available memory. The failure arises from a mismatch between the model workload and the machine budget.

The silicon contract says that every model architecture makes a computational bargain with the hardware it runs on. On the algorithm axis, the architecture's mathematical operators determine the required work and state. Precision, implementation, data reuse, runtime, and hardware then determine the realized memory traffic, latency, and energy. To honor the contract, a systems engineer must understand both the operators and how they map to the machine.

The operators that follow are not abstract theory but a specification for computational workloads. Neural computation shifts the emphasis from explicit logical instructions (if-then-else) toward large sequences of continuous mathematical transformations (multiply-add-accumulate). This shift from *Logic* to *Arithmetic* creates dense tensor workloads whose bottleneck depends on arithmetic intensity, batch size, reuse, and the hardware roofline (the performance envelope set by a chip's peak compute and memory bandwidth). A failure in such a system may come from numerical instability, a gradient that shrinks toward zero, or an activation function that stops changing with its input rather than from syntax. Concretely, recognizing a single handwritten digit in the running MNIST network requires 109,184 MACs without a logical branch in the model computation.

Arithmetic without branches is not arithmetic without risk: a number that exceeds the range its format can represent can trigger an exception or silently become an infinite or invalid value, depending on the operation and runtime. A canonical example comes from outside machine learning entirely.



War Story 5.1: The overflow that became guidance (1996)

Context: The European Space Agency's Ariane 5 Flight 501 reused inertial reference software from Ariane 4, including numeric conversions whose assumptions matched the older vehicle's flight profile (European Space Agency 1996).

Mechanism: Ariane 5's flight profile caused a 64-bit floating-point value to exceed the range of a signed 16-bit integer during conversion, triggering an unhandled operand error.

Impact: The flight computer crashed 36 seconds after launch, causing the launcher to veer off course, break apart, and self-destruct, destroying the \$370 million satellite payload.

Response: The inquiry recommended limiting software to functions required during flight, reviewing range assumptions, and designing exception handling so a single conversion failure could not disable redundant units (European Space Agency 1996).

Systems lesson: Numerical ranges, exception handling, and reuse assumptions are part of the systems contract. A computation can be syntactically correct and still be invalid for the physical regime in which it runs. ML systems hit this whenever a numerical format's representable range fails to match the magnitudes flowing through it: a low-precision multiplication that overflows produces an Inf or NaN that propagates silently through the rest of the computation, making numerical regimes as important to debugging as control flow.

Definition 5.1: Deep learning

Deep learning is the computational paradigm that learns layered feature representations from data by composing nonlinear transformations. For suitable tasks, those learned representations can reduce manual feature engineering while increasing computation (O) and state.

1. **Significance:** For suitable compositional function classes, depth can represent a hierarchy much more parameter-efficiently than a shallow network. Each layer adds work and state, and the resulting binding constraint may lie on the data, algorithm, or machine axis depending on the workload.
2. **Distinction:** Shallow methods such as logistic regression and support vector machines can work well with appropriate representations. Deep learning instead uses multiple nonlinear stages to learn representations jointly with the prediction task, which can reduce manual feature design and sometimes support transfer across tasks.
3. **Common pitfall:** A frequent misconception is that parameter count alone explains deep learning. Width can approximate broad function classes, but depth can encode some compositional functions far more efficiently. Architecture, data, optimization, and scale together determine the practical result.

The analysis treats deep learning as a physical computation workload: examining the operations used by the feed-forward networks in this chapter, evaluating how they compose into architectural structures, and measuring what training and inference demand from hardware across the D·A·M taxonomy. A single handwritten-digit recognizer makes each cost concrete. The landmark *Nature* review by LeCun, Bengio, and Hinton¹ (LeCun et al. 2015) synthesized this paradigm.

Classical machine learning often required experts to design feature extractors for each new problem, a labor-intensive process that encoded domain knowledge into handcrafted representations. Deep learning can reduce this burden by learning representations from data through hierarchical layers of nonlinear transformations. To see where neural networks fit in the broader landscape, figure 5.1 combines two views: a historical timeline of major milestones and a taxonomic nesting in which deep learning is a subset of machine learning, which is itself a subset of artificial intelligence.

This paradigm shift adds engineering problems that conventional software debugging does not address. Deep learning failures can have subtle symptoms: gradient instabilities² that prevent learning, numerical precision errors that corrupt model weights over many iterations, or inefficient memory access patterns in tensor operations³ that leave accelerator compute units underutilized. These are systems problems that require understanding the mathematical machinery underneath, not only tracing a line of code.

Diagnosing and solving such layout mismatches requires understanding how mathematical choices translate directly into computational workloads. The structural demands of a model depend on how the system represents patterns—whether through explicit rules, engineered features, or end-to-end learned representations. A single MNIST digit makes that transition concrete as it passes through three computational paradigms, with each step reshaping the system's work.

¹ | **LeCun, Bengio, and Hinton:** Recipients of the 2018 ACM Turing Award for conceptual and engineering breakthroughs that made deep neural networks a critical component of computing. Their collective work advanced representation learning, network architectures, and training methods that underpin modern deep learning.

² | **Gradient instabilities:** In a simplified 20-layer sigmoid path, the activation-derivative factors alone are bounded above by 0.25^{20} , or 9.1×10^{-13} . The complete gradient also includes weights and other operations, but repeated saturated derivatives can still make early-layer updates negligible. Plateaus or NaNs can reveal the symptom without identifying the cause; rectified linear unit (ReLU) activations, careful initialization, normalization, optimization methods, and residual connections all helped make deep networks trainable, and section 6.8.2 treats residual connections in detail.

³ | **Tensor operations:** A tensor's logical dimensions map onto a flat physical layout, and a kernel may prefer a different ordering. If changing a typical ImageNet input ($224 \times 224 \times 3$, about 151 KB) from channel-first to channel-last order requires a materialized reorder, it moves about 301 KB of read-plus-write traffic per image without changing the model's mathematical result. Some runtimes avoid the copy through views, fusion, or a compatible kernel. When a reorder remains, it can consume part of a tight inference budget before the model's main arithmetic begins.




Figure 5.1: AI Hierarchy and Timeline: A timeline from the 1950s to the 2010s is overlaid with the present-day taxonomic relationship among artificial intelligence, machine learning, and deep learning. The nested bands show scope, not exclusive historical eras: deep learning falls within machine learning, which falls within artificial intelligence. Dated milestones run from the 1956 Dartmouth Workshop through the 2017 transformer.

5.2 Computing with Patterns

The shift from logic to arithmetic reshapes how a computer encodes real-world patterns. The analysis tracks one task across all three paradigms: classifying a handwritten digit from a 28×28 pixel image from the MNIST dataset (the same input used throughout this chapter). The computational profile changes as representation strategies evolve.

5.2.1 From explicit logic to learned patterns

Rule-based programming requires developers to specify the procedures that map inputs to outputs. Consider a simple game like Breakout.⁴ Its simulator uses explicit collision rules: when the ball hits a brick, the code removes the brick and reverses the ball's direction (figure 5.2). Such general rules work well for well-defined game physics, but hand-authoring a recognition policy for every variation in messy, unstructured real-world data does not scale.

⁴ | **Breakout (DQN - deep Q-network):** Atari's 1976 arcade game became an AI milestone when DeepMind's DQN learned a gameplay policy from pixels (2015) without a hand-authored policy for each situation. The environment, action space, reward, and preprocessing were still programmed. DQN resized frames to 84×84 , stacked 4 frames as input, and selected a new action every 4 frames, illustrating how learned control still depends on an engineered real-time inference loop.




Figure 5.2: Breakout Collision Rules: The game program uses explicit if-then rules for collision detection, specifying ball direction reversal and brick removal upon contact. While effective for a game with clear physics and limited states, this approach illustrates how rule-based systems must anticipate every possible scenario.

The data flow in figure 5.3 makes the traditional programming relationship explicit: a program applies specified procedures to input data to produce outputs. Early artificial intelligence research explored whether hand-authored symbolic rules could scale to complex reasoning and recognition tasks.

Despite their apparent simplicity, rule-based limitations surface quickly with complex real-world tasks. Recognizing human activities illustrates the challenge. Classifying movement below 4 km/h as walking seems straightforward until real-world complexity intrudes. Speed variations, transitions




Figure 5.3: Traditional Programming Flow: Rules and data serve as inputs to a traditional program, which produces answers as output. This input-output pattern formed the basis for early AI systems but lacks the adaptability needed for complex pattern recognition tasks.

between activities, and unanticipated cases each demand additional rules, as the successive code fragments in figure 5.4 show. Computer vision tasks compound these difficulties. Detecting cats requires rules about ears, whiskers, and body shapes while accounting for viewing angles, lighting, occlusions, and natural variations. Early systems achieved success only in controlled environments with well-defined constraints.



Figure 5.4: Activity Classification Rules: Successive code fragments classify walking below 4 km/h, then add running below 12 km/h and biking at higher speeds. A final unhandled activity shows how real-world cases force increasingly complex branching logic.

Recognizing these limitations, the knowledge engineering approach that characterized AI research in the 1970s and 1980s attempted to systematize rule creation. Expert systems⁵ encoded domain knowledge as explicit rules, showing promise in specific domains with well-defined parameters but struggling with tasks humans perform naturally: object recognition, speech understanding, and natural language interpretation. These failures highlighted a deeper challenge: many aspects of intelligent behavior rely on implicit knowledge that resists explicit rule-based representation.

Consider one illustrative rule-based classifier for a 28 by 28 digit: compare pixel intensities against thresholds, check stroke patterns in specific regions, and branch on the results. Under this toy implementation, the computation is roughly 100 comparisons over 784 bytes of pixel data—small, predictable, and compatible with an ordinary CPU cache. The comparison establishes a workload baseline, not the accuracy of a production rule system.

5.2.2 The feature engineering bottleneck

The failures of rule-based systems suggested an alternative: rather than encoding human knowledge as explicit rules, let the system discover patterns from data. Machine learning offered this direction—instead of writing rules for every situation, researchers wrote programs that identified patterns in examples. The success of these methods, however, still depended heavily on human insight to define which patterns to look for, a process known as feature engineering.

Feature engineering transformed raw data into representations that expose patterns to learning algorithms. The histogram of oriented gradients (HOG)⁶ (Dalal and Triggs 2005) method exemplifies this approach, identifying edges where brightness changes sharply, dividing images into cells, and measuring edge orientations within each cell (figure 5.5). This transforms raw pixels into shape descriptors robust to lighting variations and small positional changes.

Complementary methods like scale-invariant feature transform (SIFT)⁷ (Lowe 1999) and Gabor filters⁸ captured different visual patterns. SIFT detected keypoints stable across scale and orientation

⁵ | **Expert systems:** These systems convert human expertise into explicit IF-THEN rules. This knowledge-engineering approach scales poorly for tasks such as object recognition, where the relevant variation is difficult to enumerate. Even successful systems such as DEC's XCON accumulated thousands of hand-authored rules, making rule acquisition, consistency, and maintenance central engineering costs.

⁶ | **Histogram of oriented gradients (HOG):** An influential handcrafted descriptor for pedestrian detection before deep learning (Dalal and Triggs 2005). One common configuration computes gradient orientations in 8×8 pixel cells, with cell size and other choices tuned for the task. HOG's fixed computation graph can run efficiently on CPUs with predictable latency, while learned features can reduce descriptor redesign but do not automatically transfer across domains and often benefit from accelerators only as model and batch scale grow.

⁷ | **Scale-invariant feature transform (SIFT):** SIFT encodes this domain expertise in a four-stage algorithm that identifies keypoints stable across scale and rotation. The number of detected keypoints varies by image, which complicates fixed-shape batching. Systems can regularize the representation through caps, padding, pooling, or fixed encodings, but each choice adds computation or discards information.




Figure 5.5: HOG Method: A local image patch is divided into cells, gradient orientations are accumulated into a histogram, and the resulting descriptor summarizes local shape while reducing sensitivity to illumination. Diagram reproduced from Figure 5 of Bakheet and Al-Hamadi (2021), licensed under CC BY 4.0.

8 | Gabor filters: Originally developed for localized time-frequency analysis (Gabor 1946), these filters detect edges and textures at specific orientations and frequencies. A typical bank contains many hand-designed filters across orientations and frequencies. Deep convolutional layers instead learn small kernels from data; early convolutional neural network (CNN) filters often become edge-, color-, and texture-sensitive detectors, shifting feature design from manual filter banks to data-driven optimization (Krizhevsky et al. 2012; LeCun et al. 2015).

changes, while Gabor filters identified textures and frequencies. Each encoded domain expertise about visual pattern recognition.

These engineering efforts enabled advances in computer vision during the 2000s. Systems could now recognize objects with some robustness to real-world variations, leading to applications in face detection, pedestrian detection, and object recognition. Despite these successes, the approach had limitations. Experts needed to carefully design feature extractors for each new problem, and the resulting features might miss important patterns that were not anticipated in their design. The bottleneck remained: human expertise could not scale to the complexity and diversity of real-world visual patterns.

Return to the same 28 by 28 digit. In this illustrative implementation, HOG divides the image into a 7 by 7 grid of 4 by 4 cells, computes gradient magnitudes and orientations, bins them into 9 orientation histograms per cell, and produces a 441-element feature vector. A linear Support Vector Machine (SVM) classifier then performs ten dot products over that vector. The result is roughly 8,000 operations and about 2 KB of working memory—about 80× the work of the toy rule-based baseline, but still structured and well served by CPU vector units using **single instruction, multiple data (SIMD)**. Descriptor size, class count, and implementation all change this cost.

5.2.3 Automatic pattern discovery

The limitations of handcrafted features motivate a more radical approach: rather than encoding features by hand, the system discovers its own. Neural networks embody exactly this shift—rather than following explicit rules or relying on human-designed feature extractors, the system learns representations directly from raw data.

Supervised learning changes the traditional programming relationship. Instead of encoding every decision rule directly, engineers provide examples and target answers, and optimization produces a learned model. Figure 5.6 makes this relationship tangible by showing learned rules as the output rather than the input. Humans still shape the task through data, objectives, architectures, and evaluation.




Figure 5.6: Data-Driven Rule Discovery: The flow diagram inverts the traditional programming pattern: data and answers serve as inputs to the machine learning process, which produces a learned decision function as output. This shift replaces a hand-coded decision policy while leaving engineers responsible for the data, objective, architecture, and evaluation.

The system discovers patterns from examples through this automated process. When shown millions of images of cats, it learns to identify increasingly complex visual patterns, from simple edges to combinations that constitute cat-like features. This parallels how biological visual systems operate, building understanding from basic visual elements to complex objects.

The gradual layering of patterns reveals *why* neural network depth matters. For some compositional function classes, deeper networks can represent functions that require much wider shallow networks, an advantage formalized in section 5.3.1.

Deep learning often exhibits empirical scaling: within a given regime, performance can improve as additional data, compute, and model capacity are used effectively, although gains depend on data quality, optimization, architecture, and evaluation conditions. In the ImageNet competition, traditional methods achieved approximately 25.8 percent top-5 error in 2011. AlexNet⁹ reduced this to 15.3 percent in 2012. By 2015, ResNet achieved 3.6 percent top-5 error, lower than a published human baseline of approximately 5.1 percent under that evaluation protocol.

Figure 5.7 previews a phenomenon observed under some datasets, model families, and training procedures. The underlying mechanisms (training error, overfitting, gradient-based learning) are developed in subsequent sections; this section establishes the possible shape. The classical regime follows a U-shaped error curve, the interpolation threshold marks sufficient capacity to fit the training sample, and error can decrease again in an overparameterized regime. The axes are normalized to emphasize shape rather than a specific dataset.




Figure 5.7: The Double Descent Phenomenon: An illustrative normalized curve shows test error first decreasing and then rising in the classical regime, peaking at the interpolation threshold, and decreasing again in the overparameterized modern regime.

The counterintuitive shape matters because test error, the error on held-out examples rather than the training examples the model sees directly, initially follows the expected U-curve, then decreases again when the model has more parameters than the simplest theory expects. This behavior complicates the classical account of overparameterization. The usual bias-variance trade-off¹⁰ suggests that models that are too small underfit, while models with excessive capacity risk fitting noise. Double descent (Belkin et al. 2019) shows that larger models can sometimes generalize better than smaller ones after the interpolation threshold. This overparameterization effect means that scale can be an engineering lever, not that scale is automatically safe: the data distribution, training procedure, and regularization still determine whether the extra capacity helps or memorizes. Overfitting and the regularization techniques that control it receive formal treatment in section 5.4.4.8; this preview needs only the shape of the curve.

Neural network performance follows empirical scaling relationships with direct systems consequences. The durable anchor is that frontier model sizes and training compute budgets have grown by orders of magnitude over the past decade (section 5.2.7 quantifies the trajectory). That growth makes arithmetic throughput, memory bandwidth, storage capacity, and communication potential bottlenecks whose importance depends on the workload and machine. Chapter 8 develops the

⁹ | **AlexNet's memory split:** Krizhevsky's team split AlexNet across two NVIDIA GTX 580 GPUs, each with 3 GB of VRAM, so that the model and training state fit across the two devices. The workaround illustrates the systems lesson: a training configuration can exceed one device's memory budget, forcing engineers to partition work and communication across devices.

¹⁰ | **Bias-variance trade-off:** In some overparameterized regimes, test error decreases again after the interpolation threshold, so the classical single-U-shaped account is incomplete. The systems consequence is that model size cannot be treated as a universally monotonic overfitting risk. Larger models can become useful engineering options, provided the measured data, optimization, and regularization regime supports the gain.

quantitative scaling formulations, including how model size, data, and compute trade off against one another; chapter 10 explores the practical responses.

Neither relationship makes model size a decision rule. Double descent describes measured generalization under a particular data and training regime, while historical scaling describes resource investment across selected systems. Choosing a model still requires comparing candidates on representative deployment data and on the complete cost of training, serving, monitoring, and maintenance. The useful question is not whether additional scale can help, but whether its measured gain earns those system costs. That comparison should also include uncertainty in both gains and costs.

🔍 Systems Perspective 5.1: When to use neural networks

Not every problem benefits from deep learning. Neural networks are strong candidates when data exhibits spatial, sequential, or hierarchical structure that an architecture can exploit, simpler baselines miss the required task quality, and measured gains justify the complete system cost (table 5.1). With limited labels, low-dimensional tabular data, approximately linear relationships, hard latency limits, or strict auditability requirements, classical methods may match the useful performance of a neural network with less training, serving, and maintenance effort (table 5.2).

Table 5.1: Neural-Network Candidates: Conditions under which deep learning is likely to outperform classical methods.

Condition	Evidence to seek	Rationale
Representative data	Training and validation cover deployment conditions	Capacity cannot replace missing coverage
Exploitable structure	Spatial, sequential, or hierarchical patterns	Architecture can encode a useful prior
Baseline gap	Simpler methods miss the required task-quality target	Complexity buys a measured gain
Operational fit	Gain survives latency, memory, and cost limits	Offline accuracy alone is insufficient
Maintenance advantage	Learned features reduce recurring manual updates	Full-system maintenance becomes simpler

Table 5.2: Classical-Method Candidates: Conditions under which simpler models are likely to match or beat a neural network.

Condition	Baseline to test	Reason to prefer it
Limited labeled data	Linear model or tree ensemble	Lower variance and easier validation
Structured tabular inputs	Gradient boosting or linear model	Strong accuracy without learned features
Approximately linear signal	Linear or additive model	Transparent and inexpensive execution
Hard resource bound	Rules or compact classical model	Predictable latency, memory use, and energy demand
Direct auditability needed	Constrained tree or linear model	Decisions are easier to inspect

Systems insight: Establish a simple baseline before building a neural network. Compare candidates on the task metric that matters and on the resources required to train, serve, monitor, and update them. Retain the neural approach only when its measured benefit justifies that added system. The document-recognition history in section 5.6 shows why the comparison must include preprocessing, rejection policy, and the cost of human review rather than model accuracy alone.

Learning representations from data reshapes AI system construction. Reducing manual feature engineering introduces new demands: infrastructure to handle large datasets, high-throughput hardware to process them, and specialized accelerators to perform mathematical calculations efficiently. These computational requirements have helped drive the development of chips optimized

for neural network operations. Deep learning now underpins major systems in computer vision, speech recognition, game playing, and natural language processing.

Return to the same 28 by 28 digit, now processed by the chapter’s three-layer neural network (784 → 128 → 64 → 10). Under the stated implementations, the forward pass requires 109,184 MACs, 1,091.8× the work of the toy rule-based baseline. The 109,386 parameters consume 438 KB in 32-bit single-precision floating-point (FP32), exceeding most L1 caches and creating traffic between cache levels during inference. Training multiplies the cost further: each image is processed forward and backward before an update, with simplified dense-layer arithmetic on the order of three forward passes, repeated over 60,000 images for multiple epochs. Dense matrix multiplications now dominate this example’s arithmetic and favor parallel hardware. This workload change raises a practical question about when engineers should invest in neural networks rather than simpler alternatives.

5.2.4 Computational infrastructure requirements

The MNIST running example traced one set of illustrative implementations from ~100 comparisons (rule-based) through ~8,000 operations (HOG) to 109,184 MACs (neural network): a 1,091.8× escalation in the stated operation models, with a corresponding shift toward parallel matrix operations. Table 5.3 summarizes the tendencies exposed by this comparison.

The comparison matters because each step changes the likely bottleneck: from branch-heavy CPU control, to batch-oriented feature pipelines, to memory-fed matrix parallelism. CPUs remain effective for small networks and latency-sensitive workloads, while large dense operations often map more efficiently to accelerators.


The shift toward parallelism creates new bottlenecks. The central challenge is the memory wall:¹¹ compute throughput has grown faster than off-chip memory bandwidth. Whether a matrix operation is memory or compute bound depends on its arithmetic intensity: large, well-tiled matrix-matrix multiplication can be compute bound, whereas matrix-vector, small, skinny, or low-reuse layers are often limited by bandwidth.¹² Hardware responses are examined in section 11.5.1, while section D.3.3 details the memory hierarchy.

Table 5.3: System Resource Tendencies: The illustrated implementations shift from control-oriented rules to structured feature pipelines and then to matrix-heavy neural computation. These are workload tendencies rather than requirements of every program in each paradigm.

System Aspect	Traditional Programming	ML with Features	Deep Learning
Computation	Often control-flow oriented	Structured parallel ops	Often matrix-parallel
Memory Access	Workload dependent	Commonly batch-oriented	Large tensor working sets
Data Movement	Program and input dependent	Feature-pipeline movement	Parameter and activation movement
Hardware Needs	Commonly CPU-centric	Often CPU with vector units	Frequently accelerator-friendly
Resource Scaling	Input and program dependent	Descriptor and data driven	Driven by width, batch, state, and reuse

Energy adds a constraint alongside speed. Moving data from main memory to processing units can consume far more energy than a simple arithmetic operation (Horowitz 2014). This energy hierarchy explains why neural network accelerators maximize data reuse: they keep frequently accessed values in local storage and schedule operations to reduce movement. Whether movement or arithmetic dominates depends on precision, reuse, workload shape, and hardware, but the cost of movement motivates specialized architectures from data center GPUs to TinyML accelerators.

The memory-computation trade-off manifests differently across the cloud-to-edge spectrum introduced in chapter 2. Cloud servers may use more memory and power to pursue throughput, while mobile devices operate within tighter power budgets. Large-batch training commonly emphasizes throughput; latency-sensitive or battery-powered inference may emphasize response time or energy per prediction. The exact priority depends on the deployment workload.



In this MNIST comparison, learned features spend far more arithmetic.

11 | **Memory wall:** Fast on-chip storage responds in nanoseconds, while larger off-chip memory is orders of magnitude farther away in latency and energy. Neural network weights often exceed the smallest caches, so performance depends on tiling, larger cache levels, on-chip static RAM (SRAM), and high-bandwidth memory (HBM) to reuse values before another off-chip fetch. Insufficient reuse can leave compute units idle even when peak arithmetic throughput is high.

12 | **Memory-bound operations:** Matrix multiplication’s arithmetic intensity (FLOP/byte loaded) determines whether a layer is compute bound or memory bound. Layers that fall below the hardware’s roofline crossover point finish their arithmetic before the next tile of weights arrives from memory. The result: effective hardware utilization can drop sharply, and adding more compute units yields little speedup until memory bandwidth or data reuse improves.

These single-machine constraints compound when scaling across multiple machines: dense layers scale with adjacent dimensions, batches scale activation storage and throughput demand, and training adds backward-pass, activation, and optimizer-state multipliers. Model-compression, hardware-acceleration, and training-system techniques all respond to this pressure by reducing the work, raising the useful machine rate, or changing where state lives.

The infrastructure demands traced earlier (parallelism, memory walls, and energy-dominated data movement) arise from how neural networks compute: weights that change during training, many simple units operating simultaneously, layers that compose low-level features into high-level concepts, and data reuse that minimizes energy-intensive movement. These properties manifest concretely in the fundamental building block of neural computation: the artificial neuron.¹³ Just as understanding a single transistor reveals how complex processors work, understanding the artificial neuron reveals how million-parameter networks operate.

5.2.5 The artificial neuron as a computing primitive

The basic unit of neural computation, the **artificial neuron** (or node), serves as a simplified mathematical abstraction of nervous activity (McCulloch and Pitts 1943). Later digital neural-network implementations adapted this abstraction into a standardized processing unit. This building block enables complex networks to emerge from simple components working together. Compare the biological and artificial neurons side by side in figure 5.8 to see how this computational model distills biological complexity into a simpler computational form.




Figure 5.8: Biological-to-Artificial Neuron Mapping: Side-by-side comparison showing how biological neuron structures map to artificial neuron components. Dendrites correspond to inputs, synapses to weights, the cell body to the summation function, and the axon to the activation output. Read the mapping as an abstraction rather than a resemblance claim, since it keeps only what reduces to arithmetic and discards the electrochemistry that produced the behavior.

The right panel of figure 5.8 traces the signal through four stages, each translating a biological structure into a mathematical operation:

1. **Input Reception** (Dendrites $\rightarrow x_1, x_2, \dots, x_n$): The artificial neuron receives a vector of input features x . In MNIST digit recognition, these represent pixel intensities. The dendrite analogy captures reception of many signals without claiming biological equivalence.
2. **Weighted Modulation** (Synapses $\rightarrow w_1, w_2, \dots, w_n$): Each input is multiplied by a learnable weight w_i , just as synaptic strengths modulate biological signals. These weights act as “gain” controls, determining how much influence each feature has on the final decision. A bias term b (shown as the top input $x_0 = 1$ in figure 5.8) shifts the activation threshold. Across a network, learned behavior is distributed over many such parameters rather than stored in one weight.
3. **Signal Aggregation** (Cell Body $\rightarrow$ Linear Function z): The neuron integrates the weighted signals, producing a single scalar value $z = \sum(x_i \cdot w_i) + b$. This mirrors how a biological cell body sums incoming electrochemical signals to determine whether the neuron has received enough evidence for a particular pattern.
4. **Nonlinear Activation** (Axon $\rightarrow$ Activation Function f): The aggregated signal passes through an activation function $f(z)$, producing the output y . This mirrors the axon’s all-or-nothing

firing decision: the nonlinearity determines whether the neuron “fires” a signal to the next layer. Unlike the biological case, f can produce graded outputs (for example, ReLU passes positive values through, zeroes negatives), but the principle is the same—thresholding followed by propagation. Table 5.4 formalizes these correspondences.

Table 5.4: **Biological vs. Artificial Neuron Mapping:** Structural components mapped to mathematical formulations and physical memory states.

Biological Structure	Artificial Component	Mathematical Operation	Engineering Role
Dendrites (receive signals)	Input Vector	$\mathbf{x} = [x_1, \dots, x_n]$	Data ingestion from sensors or prior layers
Synapses (modulate strength)	Weight Vector	$\mathbf{w} = [w_1, \dots, w_n]$	Learnable parameters encoding importance
Cell Body (integrates signals)	Linear Function z	$z = \sum (x_i \cdot w_i) + b$	Linear integration of feature signals
Axon (fires output)	Activation Function f	$a = f(z)$	Nonlinear thresholding and signal propagation

From a systems engineering perspective, this translation reveals *why* neural networks can have demanding computational requirements. A direct scalar evaluation of one neuron performs N **multiply-accumulate (MAC)**¹⁴ operations and, without reuse, loads N inputs, N weights, and a bias before writing the output. Batched and tiled kernels reuse inputs and weights across many neuron evaluations, reducing traffic per MAC. At network scale, performance therefore depends on both the number of operations and how effectively the implementation reuses data.

The transition from individual neurons to integrated systems requires navigating the central trade-off between representational capacity and computational cost. Digital hardware performs arithmetic quickly, but the volume of operations and data movement in deep networks creates bottlenecks that a single-neuron view does not expose.

Replicating intelligent behavior in silicon confronts three interrelated system-level constraints. As models grow, moving parameters and activations can become more limiting than arithmetic throughput. Concurrency clashes with dependency: many operations inside a layer can run in parallel, but layer $\ell + 1$ still depends on layer ℓ , placing a lower bound on end-to-end latency. Precision also trades against memory, throughput, and energy: wider representations consume more storage and data movement, while narrower formats require evidence that task quality remains acceptable. Chapter 10 develops this search for minimum viable precision.

Addressing these constraints requires two complementary strategies. An architectural inductive bias is a built-in structural assumption about the data: convolutional networks assume nearby pixels matter together in images, while recurrent networks assume order matters in sequences. Encoding those assumptions restricts the function class toward structures expected to be useful (Mitchell 1980). Computational scaling supplies additional capacity and optimization effort. Modern AI engineering combines both approaches, using architectural structure to make effective use of available scale.

5.2.6 Hardware and software requirements

Translating neural concepts to silicon carries a physical cost. Feature extraction can become weighted linear sums followed by nonlinear activations, while dense interactions become matrix operations that hardware must execute efficiently. The mathematical graph, however, does not prescribe an execution plan. Frameworks and compilers must choose kernels, tensor layouts, fusion boundaries, and schedules for a particular processor. Two implementations can therefore compute the same function with the same nominal arithmetic count yet differ substantially in data traffic, temporary storage, and latency. Hardware supplies arithmetic units, memory, and communication links; software determines how the graph uses them. This translation layer decides whether the parallelism visible in the model becomes useful machine utilization.

A neural network’s learned behavior is distributed across parameters stored at ordinary memory addresses. Inference fetches parameters and activations; training adds saved intermediates, gradients, and optimizer state, so storage capacity and bandwidth both matter. Biological synapses

14 | **MAC (multiply-accumulate):** A core operation in dense and convolutional layers is $a \leftarrow a + (b \times c)$. Hardware data sheets often report fused multiply-add throughput in FLOPs because one multiply and one add count as two floating-point operations. On that convention, the reported FLOP/s rate is twice the MAC/s rate when a MAC is counted as one multiply-accumulate. Layer and batch choices also create activation, communication, and data-movement costs, so MAC count is necessary but not sufficient for a systems budget.

combine storage and local processing, whereas digital accelerators move values through a memory hierarchy. The brain consumes about 20 percent of the body's resting oxygen and calories (Raichle and Gusnard 2002), although direct efficiency comparisons are misleading because biological and artificial systems perform different tasks under different accuracy and latency constraints. The narrower systems lesson is that data movement matters enough to motivate the specialized hardware in chapter 11 and the optimization strategies in chapter 10.

These hardware demands did not emerge overnight. The tension between algorithmic ambition and available silicon has shaped the entire trajectory of neural network research, from the earliest perceptrons to today's trillion-parameter models.

5.2.7 Evolution of neural network computing

The **perceptron**¹⁵ introduced a probabilistic neuron model for learning and information storage (Rosenblatt 1958). Its simple mathematics met a hardware constraint: larger networks needed far more processing and memory. Deep learning evolved through this algorithm-silicon co-evolution; the neuron abstraction remained recognizable while the hardware able to execute it transformed. That co-evolution frames the history that follows: progress required both better methods and better machines.

The **backpropagation**¹⁶ algorithm was applied to neural networks by Paul Werbos in his 1974 PhD thesis (Werbos 1974), building on Seppo Linnainmaa's 1970 work on automatic differentiation (Linnainmaa 1970), and was later popularized by Rumelhart, Hinton, and Williams (Rumelhart et al. 1986). Their publication demonstrated the algorithm's practical effectiveness and brought it to widespread attention in the machine learning community, triggering renewed interest in neural networks. This chapter returns to the algorithm in section 5.4.4; section 8.3.3 develops its systems-level implementation. The 1986 demonstration showed that hidden units could learn useful internal features, but it did not resolve the systems problem of scaling training. Larger models multiply forward operations, saved activations, gradient computations, and parameter updates. Practical adoption therefore depended on improvements across the complete stack rather than on the learning rule alone.

This history yields a recurring systems lesson: an algorithm becomes useful only when the complete system can train and evaluate it effectively. The lag between early backpropagation work¹⁷ and widespread adoption reflected the maturation of hardware, data, initialization, nonlinearities, optimization, benchmarks, and software infrastructure. An algorithm can therefore be viable in principle long before its surrounding system makes it useful in practice. Deep learning came from this convergence, not from a new mathematical discovery alone.

The term itself gained prominence in the 2010s, coinciding with advances in computational power and data accessibility. The scale of this computational explosion is difficult to grasp without visualization. Figure 5.9 plots representative estimates of AI training compute across nearly seven decades on a logarithmic scale, revealing two fitted regimes: total training compute, measured in floating-point operations (FLOPs), followed a comparatively slow pre-2010 trend, then the post-2012 deep-learning frontier accelerated sharply. Large-scale models after 2015 sit orders of magnitude above the pre-2010 trajectory, showing that modern progress reinvests hardware and algorithmic gains into much larger training runs.

Beyond raw compute, this growth carries an energy cost that systems engineers cannot ignore. Under the stated scenario assumptions, a three-day LeNet-1 workstation run would use roughly 54 kWh. Applying an external GPT-4 GPU-day estimate and a data-center-overhead factor yields about 33,600 MWh, comparable to the annual electricity use of roughly 3,200 representative US homes.¹⁸ These are order-of-magnitude estimates, not measured utility records, and they vary with accelerator utilization, facility overhead, and the calculation boundary. Even so, the scale makes energy per operation and achieved utilization design concerns alongside peak FLOP/s. The quantitative energy analysis appears in chapter 11.

Table 5.5 grounds these trends in representative systems, showing how parameters, compute, and hardware co-evolved across four decades of neural network development. Three patterns deserve separate treatment. First, the plotted post-2012 training-compute frontier rises on a timescale of months, although the fitted rate depends on the selected systems and on uncertain estimates. Second, algorithmic and systems advances have reduced the compute needed to reach fixed benchmark

¹⁵ | **Perceptron:** A machine built to execute a learning algorithm, directly linking hardware and software from the start. A single linear-threshold layer cannot represent nonlinearly separable functions such as XOR, regardless of additional input hardware. Multilayer nonlinear networks expand the representable function class, but they also introduce the training and systems costs developed in this chapter.

¹⁶ | **Backpropagation:** Short for "backward propagation of errors," the algorithm uses the chain rule to compute how the loss is locally sensitive to millions of weights. It does not establish which weights historically caused an error. Werbos applied it to neural networks in 1974 (Werbos 1974), and the 1986 Rumelhart, Hinton, and Williams publication demonstrated practical effectiveness (Rumelhart et al. 1986). The systems cost: backprop requires storing forward-pass activations and additional training state, creating the several-times-higher memory footprint quantified later in the chapter.

¹⁷ | **Algorithm-system adoption lag:** Backpropagation was available in neural-network form by 1974 (Werbos 1974), while the 1986 Rumelhart, Hinton, and Williams demonstration brought it much wider attention (Rumelhart et al. 1986). Attention mechanisms appeared in neural machine translation in 2014, transformers made them central in 2017, and later TPU- and GPU-scale infrastructure enabled much larger deployments. These histories do not imply that every computationally expensive idea will succeed. They show that an algorithm should be evaluated together with the data, optimization methods, software, hardware, and task regime needed to make it useful.




Figure 5.9: Computational Growth: Log-scale scatter plot of representative training-compute estimates in FLOPs from 1958 to 2024. Fitted trend lines compare a slower pre-2010 regime with a steeper post-2012 deep-learning frontier; the values illustrate order-of-magnitude growth rather than an audited census of model training runs.

targets in some domains, so a frontier-compute trend does not describe efficiency at a constant capability. Third, hardware utilization, reduced precision, software maturity, and procurement conditions mediate the relationship between raw compute and monetary cost; they do not guarantee that cost grows more slowly. For systems engineers, these distinctions affect infrastructure timelines, experimental budgets, and build-versus-buy decisions. Planning therefore requires measured efficiency, a stated comparison target, and explicit uncertainty, especially for closed models whose architecture and training configuration are undisclosed.

Table 5.5: Historical Performance: Four decades of representative neural-network systems showing growth in model state, training compute, and hardware infrastructure. Earlier rows use published specifications where available. GPT-3’s hardware configuration and duration are undisclosed; GPT-4’s architecture, parameter count, hardware configuration, and duration are undisclosed, and its compute entry is an external order-of-magnitude estimate consistent with figure 5.9. Total parameter count also does not determine per-input computation for sparsely activated architectures. “OoM” denotes order-of-magnitude uncertainty.

Year	System	Params	Train FLOPs	Hardware	Train Time	Error/Task
1989	LeNet-1	~10K	10^{11} – 10^{12}	Sun-4/260 workstation	3 days	1% at 12.1% rejection (USPS digits)
1998	LeNet-5	$60\text{K} \pm 1\text{K}$	$10^{14} \pm 1$ OoM	SGI Origin 2000 (200 MHz)	2–3 days	0.95% (MNIST)
2012	AlexNet	~60M	5×10^{17}	2× GTX 580 GPUs	5–6 days	15.3% (ImageNet)
2015	ResNet-152	~60M	$10^{19} \pm 0.5$ OoM	8× Tesla K80 GPUs	~3 weeks	3.6% (ImageNet)
2020	GPT-3	175B (exact)	3×10^{23}	Undisclosed	Undisclosed	N/A (language)
2023	GPT-4	Undisclosed	10^{24} – 10^{25} (external est.)	Undisclosed	Undisclosed	N/A (language)

Advances across three dimensions reinforced these trends: data availability, algorithmic innovations, and computing infrastructure. Follow the arrows in figure 5.10 to see the interaction: faster infrastructure made larger data and models practical, growing datasets created new opportunities and pressures for algorithms, and algorithmic advances changed the demands placed on computing systems. The cycle is enabling rather than automatic: progress in one dimension can expose a bottleneck in another rather than produce an immediate gain.

18 | Training energy scale: The estimate assumes a 750 W workstation for the historical run and uses 10.5 MWh/year as a representative US household electricity budget. The closed-model scenario treats an external GPU-day estimate as A100-equivalent accelerator time and applies a simple data-center overhead. The result illustrates scale; it does not establish the actual energy used to train GPT-4.




Figure 5.10: Deep Learning Virtuous Cycle: Three mutually reinforcing factors, data availability, algorithmic innovations, and computing infrastructure, form a self-reinforcing loop where breakthroughs in one area create opportunities in the others.

Data availability supplied the volume of examples that learned representations require. The rise of the internet and digital devices created vast new sources of training data: image sharing platforms provided millions of labeled images, digital text collections enabled language processing at scale, and sensor networks generated continuous streams of real-world data. This abundance provided the raw material neural networks needed to learn complex patterns effectively.

Algorithmic innovations turned that volume into trainable systems. New methods for initializing networks and controlling learning rates made training more stable. Techniques for controlling overfitting improved generalization in measured regimes. Researchers also documented empirical relationships among model size, computation, data quantity, and performance, leading to increasingly ambitious architectures while leaving the useful scaling regime task dependent.

GPT-4 33600 MWh

LeNet 54 kWh

log scale

Neural-network history is a scale explosion in training energy.



Checkpoint 5.1: Understanding deep learning's emergence

Before connecting this history to its infrastructure consequences, verify your understanding of why deep learning emerged:

- ☐ **Rule-based limits:** can you explain why rule-based programming fails for tasks like image recognition?
- ☐ **Classical ML vs. deep learning:** can you compare classical ML's hand-engineered features with deep learning's automatic feature learning?
- ☐ **Three enabling factors:** can you describe the three factors that converged to enable modern deep learning (data, algorithms, infrastructure)?
- ☐ **Hardware demand:** can you explain why deep learning demands specialized hardware more often than traditional programming?

If any of these concepts remain unclear, review the relevant sections before continuing. The systems consequences that follow build directly on this conceptual foundation.

The resulting workloads created demand for higher-throughput computing infrastructure, which evolved in response. On the hardware side, GPUs provided the parallel processing capabilities needed for efficient neural network computation, and Tensor Processing Units (TPUs)¹⁹ (Jouppi et al. 2023) pushed performance further. High-bandwidth memory systems and fast interconnects addressed data movement challenges. Software advances matched the hardware evolution: frameworks and libraries simplified building and training networks, distributed computing systems enabled training at scale, and tools for optimizing model deployment reduced the gap between research and production.

The convergence of data availability, algorithmic innovation, and computational infrastructure created the foundation for modern deep learning. The preceding checkpoint consolidates this arc before the chapter returns to the computational operations that drive it.

The historical trajectory from perceptrons through periods of reduced interest to GPU-era deep learning reveals a recurring systems pattern: an algorithm may exist before data, optimization,

¹⁹ | **TPU (tensor processing unit):** Google's custom accelerator, first deployed internally in 2015, was optimized specifically for the matrix multiplications that dominate neural network workloads. The TPU v1 achieved 92 TOPS (vendor-reported INT8 tera-operations/s) for inference at 75 W, a power-efficiency point that general-purpose GPUs of the era could not match. The name "Tensor Processing Unit" reflects the design decision to sacrifice general-purpose flexibility for maximum throughput on the operation neural networks need most.

software, benchmarks, and hardware make it practical. Modern frontier models again press against memory and energy budgets. Understanding the mathematical operations that create those pressures requires examining the computational primitives themselves.

5.3 Neural Network Fundamentals

The question now is *why* the computational demands can become so extreme. For suitable neural workloads, a GPU can outperform a CPU because it provides high parallel throughput and data reuse for matrix and tensor operations, not simply because of clock speed. The structure of these network primitives—neurons, layers, and nonlinear activations—determines how data flows through hardware and how memory demands scale during training and inference. Understanding these operations reveals how simple arithmetic on individual neurons compounds into the infrastructure requirements that shaped modern AI.

The concepts here apply to the feed-forward neural networks that anchor this chapter and recur in many larger architectures. Their common fundamentals are weighted sums, nonlinear activations, and gradient-based learning. Mastering these operations and their computational characteristics provides a basis for reasoning about more specialized networks, while later chapters introduce operators and execution patterns that this simple model does not contain.

5.3.1 Why depth matters: The power of hierarchical representations

A single-layer network attempting to classify handwritten digits must map raw pixels directly to labels. A deeper network can instead exploit hierarchical structure when the target function is compositional. The question is *why* depth can provide such representational advantages, and the answer grounds all the mathematical development that follows.

Deep networks can exploit **compositionality**: complex patterns may decompose into simpler patterns that themselves decompose further. In image recognition, learned representations can progress from local edges and textures toward parts and objects. Such hierarchical regularities occur in many datasets and motivate layered models, although a trained network need not recover this exact sequence.

Consider recognizing the digit “seven” in the MNIST example. A single-layer network maps all 784 pixel values directly to a decision. When the task has compositional structure, a deep network can compute the decision more parameter-efficiently by reusing intermediate features:

- **Layer 1** learns simple edge detectors: vertical lines, horizontal lines, diagonal strokes
- **Layer 2** combines edges into shapes: the horizontal top stroke of a “seven,” the diagonal downstroke
- **Layer 3** combines shapes into complete digit patterns

Each layer builds on the previous, allowing later computations to reuse earlier features. This hierarchy can be much more parameter-efficient than a shallow representation when the task has suitable compositional structure. The same edge detectors learned for “seven” can also contribute to recognizing “one,” “four,” and other digits. This **parameter reuse** helps explain why depth is an effective design choice without implying a fixed parameter advantage for every task. However, the choice between adding layers and widening existing ones is not symmetric: depth and width contribute to representational capacity through different mechanisms.

Layered processing in biological vision helped motivate hierarchical models, without implying that artificial networks reproduce the same mechanisms. The architectures examined in chapter 6 encode different structural assumptions for images, sequences, and other data types. Depth explains why layered representations can be useful; the remaining mechanics explain how the hierarchy is implemented. The following sections develop those mechanics: how neurons compute, how layers connect, and how information flows from input to output.

Systems Perspective 5.2: The depth vs. width trade-off

For certain function classes, depth provides an exponential representational advantage: a network with N_L layers can represent functions that require exponentially more units in a

shallow network (Telgarsky 2016). This restricted theoretical result explains why nonlinear composition can be parameter-efficient; it is not a universal guarantee for every task.

However, depth introduces three engineering challenges with each additional layer:

- Adds sequential dependencies (layer $\ell + 1$ waits for layer ℓ), limiting parallelism
- Increases gradient path length, risking vanishing/exploding gradients
- Requires storing intermediate activations for backpropagation

Modern architectures balance depth against width. A network with ten layers of 100 neurons has the same 1,000 total hidden units as one with two layers of 500, but not the same parameter count, function class, or execution graph. Both expose parallelism within each layer; the deeper network adds sequential dependencies and may represent some compositional functions more efficiently.

5.3.2 Network architecture fundamentals

A neural network’s architecture determines how information flows from input to output. Modern networks can be enormously complex, but they all build on a few organizational principles that shape both implementation and the computational infrastructure they demand.

Handwritten digit recognition grounds these concepts in a concrete example, specifically the task of classifying images from the MNIST dataset (LeCun et al. 1998). This seemingly simple task reveals all the core principles of neural networks while providing intuition for more complex applications.

Each architectural choice, from how neurons are connected to how layers are organized, creates specific computational patterns that must be efficiently mapped to hardware. This mapping between network architecture and computational requirements is essential for building scalable ML systems.



Example 5.1: Running example: MNIST digit recognition

Scenario: Deploying a basic handwritten digit classifier on 28×28 grayscale images across 10 output classes.

Diagnosis: The chapter’s running architecture maps 784 raw pixel features through 128 and 64 hidden units to 10 outputs, creating roughly 100K parameters.

Systems lesson: Inspecting input vector dimensions and layer parameters on a compact dataset establishes baseline compute and memory bandwidth requirements before scaling to high-dimensional production workloads.

5.3.2.1 The perceptron’s weighted sum

The computational machinery within each dense layer is a perceptron-style artificial neuron whose signal path (inputs, weights, bias, aggregation, activation) section 5.2.5 traced through its biological origins. What remains is to formalize that path mathematically, because this repeated operation becomes the layer-level kernel that hardware executes. In the MNIST network, a first-layer unit combines all 784 pixel intensities into one activation that may respond to a learned pattern, such as a vertical edge shared by “one” and “seven.” Follow figure 5.11 from left to right: each input x_i is multiplied by its corresponding weight w_{ij} , the products are accumulated with a bias, and an activation function transforms the resulting score. Different neurons repeat the same structure with different learned weights, allowing one layer to test many patterns against the same input.

Each input x_i has a corresponding weight w_{ij} , and the perceptron multiplies each input by its matching weight. The intermediate output, z , is computed as the weighted sum of inputs in equation 5.1:

$$z = \sum (x_i \cdot w_{ij}) \quad (5.1)$$

Here, j identifies the receiving neuron once perceptrons are assembled into a layer, and the sum runs over all input indices i . In plain terms, each input feature is scaled by how important it is (its weight), and the results are summed into a single score. This is the dot product of two vectors—the fundamental operation that hardware accelerators are designed to execute at maximum




Figure 5.11: Perceptron Architecture: Inputs enter on the left, each scaled by its own weight, and the weighted products plus a bias meet at the summation node before a single activation function produces the output. The multiplies and the accumulation along this path are the per-neuron unit of work that the chapter’s layer-cost calculations multiply out.

throughput, and the reason neural network performance is measured by the number of MAC operations completed per second.

A bias term b shifts the linear output up or down, giving the model additional flexibility to fit the data. Thus, the intermediate linear combination computed by the perceptron including the bias becomes equation 5.2:

$$z = \sum (x_i \cdot w_{ij}) + b \quad (5.2)$$

With this weighted sum in hand, a suitable output activation and loss can support regression or classification. A binary linear-threshold classifier selects one of two classes around a threshold; multiclass networks instead produce several scores and commonly normalize them with softmax.

Under the direct scalar, no-reuse accounting established earlier, one neuron performs N multiply-accumulate operations and moves its inputs, weights, bias, and output. What the formalization adds is the layer view. A dense layer of M neurons repeats the weighted sum M times, so its arithmetic cost is $M \times N$ MACs—exactly the matrix multiplication $\mathbf{xW}$ that hardware must execute, and the form in which the rest of this chapter counts work.

5.3.2.2 Nonlinear activation functions

Activation functions are where expressiveness, gradient flow, and hardware cost meet. They convert linear weighted sums into nonlinear outputs; without them, multiple linear layers would collapse into a single linear transformation, severely limiting the network’s expressive power. Figure 5.12 compares three commonly used element-wise activation functions and one vector-level function (softmax), each with mathematical characteristics that shape both learning behavior and execution cost.

The choice of activation function affects both learning and computational efficiency, revealing how systems constraints shape algorithmic design. ReLU ($\max(0, x)$) became a common hidden-layer activation for feed-forward networks because it is inexpensive, has local derivative one for positive inputs, and can produce sparse activations. Earlier differentiable deep networks often used sigmoid or tanh. Their smooth S-curves can saturate, causing gradients to shrink across many layers and impeding early-layer learning. That contrast is not a simple ranking: sigmoid remains appropriate for some binary outputs, tanh remains useful when bounded zero-centered activations are desired, and modern architectures use additional activation families. The relevant design question is which function supplies the required output semantics and gradient behavior at an acceptable execution cost. Understanding those roles explains ReLU’s adoption without turning it into a universal default.

Sigmoid The logistic sigmoid function²⁰ squashes continuous inputs onto the open interval $(0, 1)$ (equation 5.3):

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad (5.3)$$

²⁰ | **Sigmoid:** From Greek *sigma* + *eidos* (“sigma-shaped”), referring to the S-curve that maps inputs to the bounded $(0, 1)$ range. The mapping requires an exponential approximation, commonly implemented through specialized instructions, lookup tables, or polynomial methods. It is generally more expensive than ReLU’s comparison-and-selection operation, although the measured gap depends on precision, vectorization, and hardware. ReLU’s adoption reflected both easier gradient flow on its positive branch and cheaper execution.




Figure 5.12: Common Activation Functions: Four nonlinear activation functions plotted with their output ranges. Sigmoid maps inputs to $(0, 1)$ with smooth gradients, tanh provides zero-centered outputs in $(-1, 1)$, and ReLU outputs zero for negative inputs. Softmax (bottom-right) shows one component of a 3-element vector while the other two logits are fixed at 13 and 0. Its magnified y-axis (roughly 0 to 0.05) is not comparable to the other panels, and it shows only one component of a vector whose outputs are interdependent and sum to 1.

The S-shaped curve maps logits to $(0, 1)$, so sigmoid can parameterize a binary probability. It approaches one for large positive inputs and zero for large negative inputs. Although sigmoid is smooth and differentiable everywhere, gradient-based learning need not be: ReLU uses an assigned subgradient at zero.

Sigmoid has a significant limitation: for inputs with large absolute values (far from zero), the gradient becomes extremely small, a phenomenon called the **vanishing gradient problem**.²¹ During backpropagation, these small gradients multiply together across layers, causing gradients in early layers to become exponentially tiny. This effectively prevents learning in deep networks, as weight updates become negligible.

Sigmoid outputs are not zero-centered. When the inputs feeding a parameter update are also nonnegative, their weight gradients can become correlated in sign, producing less direct optimization trajectories. The complete gradient still depends on upstream errors and is not universally one-signed.

Tanh The hyperbolic tangent²² centers output activations around zero over the interval $(-1, 1)$ (equation 5.4):

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \quad (5.4)$$

Tanh produces an S-shaped curve similar to sigmoid but centered at zero: negative inputs map to negative outputs and positive inputs to positive outputs. This symmetry balances gradient flow during training, often yielding faster convergence than sigmoid.

Like sigmoid, tanh is smooth and differentiable everywhere, and it still suffers from the vanishing gradient problem for inputs with large magnitudes. When the function saturates (approaches -1 or 1), gradients shrink toward zero. Despite this limitation, tanh's zero-centered outputs make it preferable to sigmoid for hidden layers in many architectures, particularly in recurrent neural networks where maintaining balanced activations across time steps is important.

²¹ | **Vanishing gradient problem:** The chain rule's multiplication of Jacobian factors across layers can cause gradients to shrink. Sigmoid's derivative is at most 0.25, so along a simplified 10-layer path the activation-derivative factors alone contribute at most $0.25^{10} \approx 10^{-6}$. The complete gradient also includes weights and other operations, but repeated saturated activation derivatives can still make early-layer updates negligible.

²² | **Tanh (hyperbolic tangent):** By centering its output range on zero, tanh can reduce the correlated-update behavior associated with all-positive activations. Its computational cost is similar to sigmoid because it also depends on exponential approximations, while its saturation behavior still permits vanishing gradients.

Both sigmoid and tanh share a critical limitation: gradient saturation at extreme input values. The search for an activation function that avoids this problem while remaining computationally efficient led to one of deep learning's most important innovations.

ReLU ReLU²³ sets negative values to zero while maintaining linear identity for positive inputs (equation 5.5):

$$\text{ReLU}(x) = \max(0, x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases} \quad (5.5)$$

ReLU's characteristic shape—a straight line for positive inputs and zero for negative inputs—provides three advantages. First, its local derivative is one for positive inputs, avoiding activation saturation on that branch; the complete network gradient can still shrink or grow through weights and other operations. Second, ReLU introduces deterministic activation sparsity by zeroing negative values; this differs from dropout's stochastic masking.²⁴ Zero values do not automatically produce a speedup, because ordinary dense kernels still process their positions unless the implementation exploits sparsity. Third, ReLU uses a comparison and selection, $\text{output} = (\text{input} > 0) ? \text{input} : 0$, rather than an exponential approximation. These are local properties of the operator, while end-to-end performance depends on how it is fused with neighboring work.

ReLU is not without drawbacks. The **dying ReLU problem**—neurons that permanently output zero and cease learning—occurs when neurons become stuck in the inactive state. If a neuron's weights evolve during training such that the preactivation $z = \mathbf{w}^T \mathbf{x} + b$ is consistently negative across all training examples, the neuron outputs zero for every input. Since ReLU's gradient is also zero for negative inputs, no gradient flows back through this neuron during backpropagation: the weights cannot update, and the neuron remains dead. This can happen with large learning rates that push weights into unfavorable regions. From a systems perspective, dead neurons represent wasted capacity: parameters that consume memory and compute during inference but contribute nothing to the output. Careful initialization (He et al. 2015), moderate learning rates, and leaky ReLU variants help mitigate dying ReLUs. Batch normalization²⁵ instead stabilizes layer inputs and permits higher learning rates (Ioffe and Szegedy 2015).

Softmax Unlike *element-wise* activation functions that operate independently on each value, softmax²⁶ is a *vector-level* function: it considers all values simultaneously to produce a probability distribution. In simple classifiers, softmax is commonly used in the output layer rather than as an element-wise hidden activation; it also normalizes learned importance-score vectors in attention architectures. The softmax function is defined in equation 5.6:

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \quad (5.6)$$

For a vector of K values (often called logits²⁷), softmax transforms them into K probabilities that sum to 1. One component of the softmax output appears in figure 5.12 (bottom-right); in practice, softmax processes entire vectors where each element's output depends on all input values.

In multiclass classifiers, softmax is commonly used in the output layer. It converts arbitrary real-valued logits into normalized predicted probabilities, but those values are not necessarily calibrated confidence estimates. The class with the highest probability becomes the predicted class, and larger logits receive disproportionately higher probabilities.

The mathematical relationship between input logits and output probabilities is differentiable, allowing gradients to flow back through softmax during training. When combined with cross-entropy loss (discussed in section 5.4.3), softmax produces particularly clean gradient expressions that guide learning effectively. Beyond their mathematical properties, the choice of activation functions has direct consequences for hardware efficiency.

23 | ReLU (rectified linear unit): Unlike the costly exponential operations in prior activation functions, ReLU's $\max(0, x)$ operation maps to a simple comparison and selection. This efficiency, together with better gradient flow for positive activations, helped make the deep architectures of the AlexNet era computationally tractable (Nair and Hinton 2010; Krizhevsky et al. 2012).

24 | Dropout: Randomly deactivating neurons during training forces a network to learn redundant representations, a regularization technique used in the AlexNet-era shift toward deep vision models (Srivastava et al. 2014; Krizhevsky et al. 2012). This creates a systems-level divergence between the computational graphs for training (stochastic) and inference (deterministic). Failing to switch from the training to the inference graph is a common bug that can silently degrade accuracy.


25 | Batch normalization systems cost: BatchNorm adds two learned parameters per feature (scale γ and shift β) and behaves differently during training and inference: training uses live mean and variance from the mini-batch, while inference uses frozen running statistics; by keeping activation distributions better scaled, it stabilizes training and can permit higher learning rates, but it also makes the layer depend on batch statistics. Small batches can produce noisy mean and variance estimates, so a mathematical stabilizer becomes a systems choice about batch size, activation memory, and training-serving parity. Later architectures and large-scale training settings introduce further variants of the same batch-statistics trade-off.

Systems Perspective 5.3: Activation functions and hardware

ReLU's widespread adoption comes from hardware efficiency, alongside its avoidance of vanishing gradients. Computing $\max(0, x)$ requires a single comparison operation, while

26 | Softmax: The name reflects its role as a “soft” or differentiable version of `argmax`, a function that must evaluate an entire vector to find its maximum value. This vector-wise operation is not used as a drop-in replacement for element-wise nonlinearities such as ReLU, but it is central whenever a model must normalize a vector of scores, including classification heads and attention layers. Direct exponentiation can overflow in 32-bit floating point. Stable implementations subtract the largest logit before exponentiating; the shift leaves every softmax probability unchanged while making the largest exponent equal to one.

27 | Logits: In multiclass neural classifiers, logits are unnormalized output scores. Score differences correspond to log probability ratios after softmax, while an individual multiclass logit is not itself a calibrated log-odds quantity. Because `argmax` over logits and `argmax` over softmax probabilities select the same class, inference can skip softmax when only the top class is needed.



An illustrative logic model shows why ReLU is generally cheaper to implement.

sigmoid and tanh require computing exponentials—operations that are much more expensive in both time and energy. The transistor-tax calculation in section 5.3.2.3 models this gap as a logic-cost ratio, and chapter 11 covers the broader benchmark and accelerator-implementation implications.

5.3.2.3 The transistor tax: Logic unit cost

The activation decision has a second cost beyond gradient behavior: silicon area. In computer architecture, we measure the **logic unit cost** in terms of transistor count and energy per operation.

A ReLU datapath can be implemented with comparison and selection logic, while sigmoid or tanh requires an approximation to a transcendental function, such as a lookup table or polynomial evaluation. The illustrative assumptions here assign 50 transistor-equivalents to ReLU and 2,500 to a high-precision exponential datapath. Actual area, latency, and energy depend on precision, approximation method, sharing, throughput target, and hardware technology.

This illustrative disparity is the **transistor tax**: under the stated assumptions, the modeled silicon “price” of sigmoid is 50× that of ReLU. The exact ratio is not a hardware constant. The durable systems lesson is that simple piecewise-linear activations are generally cheaper to implement than transcendental functions, while ReLU’s adoption also reflects its optimization behavior and gradient flow.

These transformations convert the linear input sum into a nonlinear output, yielding the generic artificial-neuron computation in equation 5.7:

$$\hat{y} = f(z) = f\left(\sum_i (x_i \cdot w_{ij}) + b\right) \quad (5.7)$$

This nonlinearity expands what multilayer networks can represent. Without it, a stack of linear layers collapses algebraically into one linear transformation, regardless of how many intermediate matrices the implementation evaluates. In figure 5.13, the left panel shows a linear boundary that cannot separate the classes, while the right shows one curved boundary that a multilayer network with nonlinear activations and suitable learned parameters can represent. The figure illustrates representational possibility rather than a guarantee that training will discover that boundary. Optimization, data coverage, and architecture still determine which function the learned parameters produce.




Figure 5.13: Linear vs. Nonlinear Decision Boundaries: Two scatter plots compare classification with and without activation functions. Without activation, a straight line fails to separate the two classes. With a nonlinear activation function applied, the network produces a curved decision boundary that correctly separates the points.

The **universal approximation theorem**²⁸ establishes that sufficiently wide networks with suitable activation functions can approximate broad classes of functions arbitrarily well. This theoretical foundation, combined with the computational and optimization characteristics of specific activation functions like ReLU and sigmoid, explains neural networks’ practical effectiveness in complex tasks. The theorem, however, states a pure existence result under the assumption of unlimited width; it says nothing about the accelerator executing the computation. Physical ML systems impose hard limits on both dimensions of the width-vs.-depth trade-off: a single hidden layer wide enough to approximate a complex function requires parameter storage and activation memory that can exhaust accelerator VRAM entirely, and a network too deep to fit its activations in memory stalls the backward pass on the same constraint. The engineering problem is therefore not whether a network of sufficient width or depth *exists* but whether it fits within the memory and bandwidth envelope of the target hardware.

5.3.2.4 Layers and connections

Individual neurons compute weighted sums, apply bias terms, and pass results through activation functions. The power of neural networks, however, comes from organizing these neurons into layers. A layer is a collection of neurons that process information in parallel. Each neuron in a layer operates independently on the same input but with its own set of weights and bias, allowing the layer to learn different features from the same input data.

In a typical neural network, three layer types form the hierarchy:

1. **Input layer:** Receives the raw data features
2. **Hidden layers:** Process and transform the data through multiple stages
3. **Output layer:** Produces the final prediction or decision

In figure 5.14, data enters at the input layer, passes through multiple hidden layers that progressively extract more abstract features, and emerges at the output layer as a prediction. Each successive layer transforms the representation, building increasingly complex features—a hierarchical processing pipeline that gives deep neural networks their ability to learn complex patterns.




Figure 5.14: Layered Network Architecture: Deep neural networks transform data through successive layers, enabling the extraction of increasingly complex features and patterns. Each layer applies nonlinear transformations to the outputs of the previous layer, ultimately mapping raw inputs to desired outputs.

As data flows through the network, it is transformed at each layer to extract meaningful patterns. The weighted summation and activation process introduced for individual neurons scales up: each layer applies these operations in parallel across all its neurons, with outputs from one layer becoming inputs to the next. This creates a hierarchical pipeline where simple features detected in early layers combine into increasingly complex patterns in deeper layers—enabling neural networks to learn sophisticated representations from raw data.

5.3.3 Parameters and connections

The learnable parameters²⁹ of neural networks, weights and biases, determine how information flows through the network and how transformations are applied to input data. Their organization directly impacts both learning capacity and computational requirements.

5.3.3.1 Weight matrices

Weights determine how strongly inputs influence neuron outputs. In dense layers, these weights organize into matrices that map efficiently to linear-algebra kernels. In a layer with n input features and m neurons, the weights form a matrix $\mathbf{W} \in \mathbb{R}^{n \times m}$, where each column represents the weights for a single neuron. A matrix-vector product processes one input, while a matrix-matrix product processes a batch of inputs and can expose more parallel work and data reuse.

Recall that for a single neuron, we computed $z = \sum_{i=1}^n (x_i \cdot w_{ij}) + b$. When we have a layer of m neurons, we could compute each neuron's output separately, but matrix operations provide a much

28 | Universal approximation theorem: Under conditions on the activation, target function, compact domain, and approximation tolerance, a sufficiently wide single-hidden-layer network can approximate the target arbitrarily well. This existence result is nonconstructive: it does not specify how to find the weights or how many neurons practical training will require. For some function classes, depth provides an exponential representational advantage over shallow networks, but this separation is not universal.

29 | Parameter memory cost: Parameter count is a misleading proxy for memory importance; some layers add small learned scale or shift parameters that barely affect byte budgets but can strongly affect model behavior when training choices change. Conversely, the bulk of parameters in dense weight matrices require extra state during training: gradients plus optimizer bookkeeping beyond the stored weights themselves. A model that fits in memory for inference may therefore require several times more memory for training, a cost quantified later in the training memory budget.

more efficient approach. Rather than computing each neuron individually, matrix multiplication enables us to compute all m outputs simultaneously (equation 5.8):

$$\mathbf{z} = \mathbf{x}\mathbf{W} + \mathbf{b} \quad (5.8)$$

This equation computes every neuron's preactivation for one input: the vector $\mathbf{x}$ multiplied by the weight matrix $\mathbf{W}$ produces all m values, and the bias vector $\mathbf{b}$ shifts each one. For a batch, the input becomes a matrix and the same layer becomes a matrix-matrix product. These kernels account for much of the work in dense feed-forward networks, but whether they dominate runtime—and whether they are compute bound or memory bound—depends on their dimensions, batch size, precision, implementation, and target hardware.

This matrix organization is more than mathematical convenience; it reflects how modern neural networks are implemented for efficiency. Each weight w_{ij} represents the strength of the connection between input feature i and neuron j in the layer.

The canonical case studied here connects each neuron in a layer to every neuron in the previous layer, forming a “dense” or “fully-connected” layer. This pattern allows each neuron to use all features from the previous layer. Fully connected layers establish the mathematical principles needed here, while alternative connectivity patterns (explored in chapter 6) exploit structure in the data and can reduce work or improve statistical efficiency.

Figure 5.15 makes the dense pattern explicit by laying out a small three-layer network with every connection labeled. Every input connects to every hidden neuron, and every hidden neuron connects to every output. Each labeled edge represents one learnable weight, making the parameter count and arithmetic structure visible. Read one destination neuron at a time: its incoming edges form one column of the preceding weight matrix, and its bias shifts the resulting weighted sum. Reading all destinations together reveals the matrix operation. The numerical labels are illustrative rather than a worked forward pass; the important relationship is the matrix shape. For layers of sizes (n_1, n_2, n_3) , the weight matrices are $\mathbf{W}^{(1)} \in \mathbb{R}^{n_1 \times n_2}$ and $\mathbf{W}^{(2)} \in \mathbb{R}^{n_2 \times n_3}$, requiring $n_1 n_2 + n_2 n_3$ weights before biases.




Figure 5.15: Fully-Connected Layers: A three-layer network with dense connections between layers, where each neuron integrates information from all neurons in the preceding layer. Labeled edges and nodes illustrate the organization of weights, biases, and activations; they are not a worked numerical forward pass.

5.3.3.2 Bias terms

Each neuron in a layer also has an associated bias term³⁰. While weights determine the relative importance of inputs, biases allow neurons to shift their activation functions. This shifting is important for learning, as it gives the network flexibility to fit more complex patterns.

For a layer with m neurons, the bias terms form a vector $\mathbf{b} \in \mathbb{R}^m$. When we compute the layer's output, this bias vector is added to the weighted sum of inputs (the same form as equation 5.8),

³⁰ | **Bias terms:** Biases add one parameter per neuron, compared with n weights for a neuron with n inputs, so they account for $1/(n+1)$ of that layer's parameters. Their model-wide share and effect on accuracy depend on layer widths, architecture, and normalization. Some modern architectures omit biases in layers followed by batch normalization because the normalization transform includes a learned shift.

where the bias terms effectively allow each neuron to have a different “threshold” for activation, making the network more expressive:

$$\mathbf{z} = \mathbf{x}\mathbf{W} + \mathbf{b}$$

The organization of weights and biases across a feed-forward network follows a systematic pattern. For each parameterized affine layer ℓ , three components define its computation; the weights and biases are learned, while the activation function is usually chosen by the designer:

- A weight matrix $\mathbf{W}^{(\ell)}$
- A bias vector $\mathbf{b}^{(\ell)}$
- An activation function $f^{(\ell)}$

This yields the complete layer computation in equation 5.9:

$$\mathbf{a}^{(\ell)} = f^{(\ell)}(\mathbf{z}^{(\ell)}) = f^{(\ell)}(\mathbf{a}^{(\ell-1)}\mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)}) \quad (5.9)$$

Here, $\mathbf{a}^{(\ell)}$ (written as $\mathbf{A}^{(\ell)}$ for batches) represents the layer’s activation output. The row-vector convention is adopted throughout: each sample is a row, and the weight matrix $\mathbf{W}^{(\ell)} \in \mathbb{R}^{n_{\ell-1} \times n_{\ell}}$ maps from the previous layer’s width to the current layer’s width. With this equation in place, the core architecture concepts are complete enough to proceed.



Checkpoint 5.2: Neural network architecture fundamentals

Before proceeding to network topology and training, verify your understanding of the foundational concepts:

Use the MNIST architecture $784 \rightarrow 128 \rightarrow 64 \rightarrow 10$ as the running check.

- ☐ **Neuron equation:** can you write the neuron equation, including weighted sum, bias, and activation function?
- ☐ **ReLU vs. sigmoid:** can you explain why ReLU is computationally cheaper than sigmoid while still providing the nonlinearity a multilayer network needs?
- ☐ **Layers to matrices:** can you map input, hidden, and output layers to the corresponding weight matrices $\mathbf{W}^{(\ell)} \in \mathbb{R}^{n \times m}$?
- ☐ **Parameter count:** can you calculate the total parameter count, including both weights and biases?
- ☐ **Width, depth, cost:** can you explain how layer width, depth, and connectivity change memory footprint and arithmetic cost?

If any of these feel unclear, review the earlier sections on neural network fundamentals, artificial neurons, and parameters before continuing. The upcoming sections on training and optimization build directly on these foundations.

5.3.4 Architecture design

Architecture design determines how individual neurons, activation functions, and weight matrices connect to solve problems that no single layer can handle. That design has direct systems consequences, because each topological choice (adding a layer, widening a hidden dimension, changing connectivity) changes the parameter count, memory footprint, and arithmetic cost.

Network topology describes how individual neurons organize into layers and connect to form complete neural networks. Building intuition begins with a simple problem that became famous in AI history.³¹

The XOR example motivates a network with a hidden layer, but real-world networks require systematic consideration of design constraints and computational scale. Recognizing handwritten digits using the MNIST (LeCun et al. 1998) dataset illustrates how problem structure determines network dimensions while hidden layer configuration remains an important design decision.

³¹ | **XOR problem:** XOR is the canonical example of why network topology matters. A single linear-threshold layer cannot represent this nonlinearly separable function, regardless of how many units it contains, whereas a small network with a nonlinear hidden layer can. The exact number of units required depends on the activation and representation conventions; the durable lesson is that adding nonlinear depth changes the class of functions the network can express.



Example 5.2: Building intuition: The XOR problem

Consider a network learning the XOR function, a classic problem that requires nonlinearity. With inputs x_1 and x_2 that can be 0 or 1, XOR outputs 1 when inputs differ and 0 when they are the same.

Setup: Two inputs $\rightarrow$ two hidden neurons $\rightarrow$ one output

Example: For inputs (1, 0):

- Hidden neuron one: $h_1 = \text{ReLU}(1 \cdot w_{11} + 0 \cdot w_{12} + b_1)$
- Hidden neuron two: $h_2 = \text{ReLU}(1 \cdot w_{21} + 0 \cdot w_{22} + b_2)$
- Output: $y = \text{sigmoid}(h_1 \cdot w_{31} + h_2 \cdot w_{32} + b_3)$

Systems insight: Hidden layers do not merely add capacity; they change the class of functions the system can represent. XOR is the small example that makes depth a structural requirement rather than a decorative design choice (Minsky and Papert 1969).

5.3.4.1 Feedforward network architecture

Applying the three-layer architecture to MNIST reveals how data characteristics and task requirements constrain network design. Compare the two panels in figure 5.16 to see this architecture from both perspectives: panel (a) schematically represents a 28×28 pixel grayscale image connected to the hidden and output layers, while panel (b) shows the same values reshaped into a 784-dimensional feature vector. Matrix multiplication can operate on multidimensional arrays, but this fully connected layer defines one weight per input feature and hidden unit, so the image axes are not preserved in its parameterization. Convolutional architectures, explored in chapter 6, instead encode spatial locality directly.

The input layer's width is directly determined by the data format. For a 28 by 28 pixel image, each pixel becomes an input feature, requiring 784 input neurons. Equivalently, 28 rows and 28 columns flatten into 784 values. We can think of this either as a 2D grid of pixels or as a flattened vector, where each value represents the intensity of one pixel.

The output layer's structure is determined by the task requirements. For digit classification, we use 10 output neurons, one for each possible digit (0–9). When presented with an image, the network produces a score for each output neuron; a softmax can then convert those logits into predicted class probabilities.

Between these fixed input and output layers, there is flexibility in designing the hidden layer topology. The choice of hidden layer structure, including the number of layers to use and their respective widths, represents one of the key design decisions in neural networks. Additional layers increase the network's depth, allowing it to learn more abstract features through successive transformations. The width of each layer provides capacity for learning different features at each level of abstraction.

Widening hidden layers from 100 to 1000 neurons increases parameters from ~89.6K to ~1.8M (~358 KB vs. ~7 MB in FP32). Whether the wider network improves accuracy requires training and evaluation; the storage increase alone already matters for mobile deployment budgets.

5.3.4.2 Layer connectivity design patterns

The fully connected architecture in section 5.3.3 maximizes flexibility by connecting every neuron to every neuron in the next layer, but connectivity itself is an engineering decision. Dense, sparse, and skip patterns trade learning flexibility against parameter count, locality, and gradient flow.

Dense connectivity represents the standard pattern where each neuron connects to every neuron in the subsequent layer. In our MNIST example, connecting our 784-dimensional input layer to a hidden layer of 128 neurons requires 100,352 weight parameters (784×128). This full connectivity enables the layer to learn arbitrary linear combinations of its inputs; nonlinear activations between dense layers extend those combinations to nonlinear relationships. The number of parameters grows as the product of adjacent layer widths, so widening one layer scales them linearly while widening both scales them quadratically.




Figure 5.16: MNIST Network Topology: Two panels show the network architecture for digit recognition. Panel (a) schematically represents a 28×28 pixel image of a digit connected through hidden layers to 10 output nodes. Panel (b) shows the same architecture with the input image flattened into a 784-element vector, illustrating how spatial data enters the network.

Sparse connectivity patterns introduce purposeful restrictions in how neurons connect between layers. Rather than maintaining all possible connections, neurons connect to only a subset of neurons in the adjacent layer. This approach draws inspiration from biological neural systems, where neurons typically form connections with a limited number of other neurons. In visual processing tasks like the MNIST example, neurons might connect only to inputs representing nearby pixels, reflecting the local nature of visual features.

As networks grow deeper, the path from input to output becomes longer, potentially complicating the learning process. Skip connections address this by adding direct paths between nonadjacent layers. These connections provide alternative routes for information flow, supplementing the standard layer-by-layer progression. In our digit recognition example, skip connections might allow later layers to reference both high-level patterns and the original pixel values directly.

These connection patterns have direct hardware consequences as well as theoretical ones. Dense connections maximize learning flexibility and map efficiently onto general matrix multiply (GEMM) kernels on tensor-oriented hardware. Sparse connections reduce theoretical parameter counts and FLOPs, but a sparse weight matrix does not automatically execute faster than a dense one on standard hardware—without structured sparsity support (such as 2:4 patterns that hardware can exploit through compressed index formats), the arithmetic reduction does not translate to proportional reductions in execution time. Skip connections help maintain effective information flow in deeper networks while preserving the gradient paths that make very deep architectures trainable.

5.3.4.3 Model size and computational complexity

How parameters (weights and biases) are arranged determines both learning capacity and computational cost—this is the model’s side of the silicon contract (section 1.7): the parameter count, their numerical precision, and the operations they require collectively define the computational bargain the model strikes with hardware. While topology defines the network’s structure, parameter initialization and organization directly affect learning dynamics and final performance.

Worked example: Training vs. inference memory A single forward pass through the $784 \rightarrow 128 \rightarrow 64 \rightarrow 10$ network costs 109,184 MACs. The memory footprint during training tells a different story. We compute the footprint for this network at batch size 32 in 32-bit (4-byte) floating-point precision, then contrast it with inference requirements, accounting for parameters, activations, gradients, and optimizer state in turn.

The first contribution is the model parameters: table 5.6 tallies the weights and biases layer by layer, totaling 109,386 parameters at 4 bytes each, occupying 437.5 KB.

Table 5.6: MNIST Parameters by Layer: Weight and bias counts for the 784-128-64-10 network. Layer 1 (Input → Hidden1) dominates the parameter budget because it multiplies the 784-dimensional input by 128 hidden units; deeper layers shrink rapidly as the representation compresses toward the 10-class output.

Layer	Weights	Biases	Total Parameters
Input → Hidden1	$784 \times 128 = 100,352$	128	100,480
Hidden1 → Hidden2	$128 \times 64 = 8,192$	64	8,256
Hidden2 → Output	$64 \times 10 = 640$	10	650
Total			109,386 parameters

Activations are the next contribution. Table 5.7 records each layer’s activation tensor and its memory cost at batch size 32.

Table 5.7: MNIST Activation Memory at Batch Size 32: Per-layer activation tensor shapes, value counts, and 32-bit float memory cost during a single forward pass. Input activations dominate the activation table at 100.4 KB; for this small MNIST multilayer perceptron (MLP), the total activation footprint remains below the parameter footprint, while larger batches and deeper networks make activations a central training-memory concern.

Layer	Activation Shape	Values	Memory
Input	32×784	25,088	100.4 KB
Hidden1	32×128	4,096	16.4 KB
Hidden2	32×64	2,048	8.2 KB
Output	32×10	320	1.3 KB
Total		31,552	126 KB

Training adds two memory contributions that inference never pays. Gradients occupy the same space as the parameters they update, here 437.5 KB, and the Adaptive Moment Estimation (Adam) optimizer stores momentum and velocity at twice the parameter size, another 875.1 KB. Table 5.8 summarizes the per-component memory footprint for training vs. inference.

Table 5.8: MNIST Training vs. Inference Memory at Batch Size 32: A simplified FP32 ledger for the toy MNIST MLP. Training retains saved activations and adds gradients plus two Adam buffers; inference uses the peak adjacent-layer buffer. Framework workspace and memory-saving techniques can change the totals.

Component	Training	Inference
Parameters	437.5 KB	437.5 KB
Activations	126 KB	117 KB
Gradients	437.5 KB	—
Optimizer state	875.1 KB	—
Total	~1.9 MB	~554 KB

Scale drives the systems lesson. At batch size 32, training requires 3.4× the inference memory in this simplified ledger. The absolute gap grows with parameter and activation storage; the ratio depends on their balance and on implementation choices such as checkpointing, precision, and workspace reuse.



Lighthouse 5.1: MNIST-to-GPT-2 memory scaling

Our MNIST running example and the GPT-2 lighthouse show how model scale changes the data (D_{vol}) term of the iron law.

- MNIST (running example): 109,386 parameters at 4 bytes each occupy approximately 438 KB. This entire model fits inside the L2 cache of a modern processor.

- GPT-2 (lighthouse): 1,500,000,000 parameters at 4 bytes each occupy approximately 6 GB before activations and runtime workspace. This exceeds cache capacity and places the model in a device-memory or multi-device regime unless precision is reduced.

Systems insight: Moving from ~109.4K to 1.5B parameters is a 13,712.9× jump. The increase represents a phase change in engineering, not merely “more parameters.” MNIST can keep its weights near the processor; GPT-2 serving must manage a much larger memory working set. Whether a particular layer is compute or memory bound still depends on batch size, precision, reuse, and hardware balance.

Parameter count grows with network width and depth. For our MNIST example, consider a network with a 784-dimensional input layer, hidden layers of 128 and 64 neurons, and a 10-neuron output layer (784 → 128 → 64 → 10). The first layer requires 100,352 weights and 128 biases, the second layer 8,192 weights and 64 biases, and the output layer 640 weights and 10 biases, totaling 109,386 parameters. Each must be stored in memory and updated during learning.

The memory requirements summarized in table 5.8 seem modest for our small MNIST classifier. Scaling to production-sized models transforms these requirements dramatically, changing the hardware regime.

Napkin Math 5.1: Quick estimation for ML engineers

Detailed calculations are essential for design documents, but experienced engineers also develop rapid mental estimation skills. These “napkin math” shortcuts enable quick feasibility checks before committing to detailed analysis.

Memory estimation:

- Parameters → bytes: Multiply by four for 32-bit single precision (FP32), two for 16-bit half precision (FP16) or brain floating point 16 (BF16, the latter with an FP32-like exponent range), or one for 8-bit integer (INT8)
- Fully connected (FC) layer parameters: Input × Output (plus Output biases, usually negligible)
- Training state: Parameters, gradients, and optimizer state can require ~3–4× the weight-only parameter memory; peak activations and workspace are additional
- Adam optimizer overhead: 2× parameter memory (momentum + velocity)
- Max batch size: Approximately (GPU VRAM – persistent state – workspace) / peak activations per sample

Compute estimation:

- FC layer FLOPs: $2 \times d_{\text{in}} \times d_{\text{out}} \times B$ (multiply-add = 2 ops)
- MACs to FLOPs: Multiply by 2
- Compute utilization proxy: achieved FLOP/s / R_{peak}

Example: “Can I train a 100M parameter model on a 16 GB GPU?”

Mental math: 100M parameters at 4 bytes each, with 4 training-memory copies, require 1.6 GB for persistent model state. That leaves about 14.4 GB for activations, temporary workspace, and framework overhead. Answer: the parameters fit, but the training configuration is not yet proven; batch size, activation shape, and kernel workspace determine whether the complete step fits.

Systems insight: A rough persistent-state estimate can reject impossible configurations quickly, but batch size still depends on activation and workspace headroom.

Beyond this model-fit example, table 5.9 distills three common feasibility questions into one-line formulas an engineer can apply before reaching for a profiler.

GPT-2 6 GB

MNIST 438 KB

log scale

MNIST is cache-scale; GPT-2 is VRAM-scale.

Table 5.9: Napkin-math feasibility checks: Three rapid mental-estimation rules for GPU fit, epoch time, and bound regime, distilled into one-line formulas an engineer can apply before running a profiler.

Question	Quick Estimate
“Will this model fit in GPU memory?”	persistent training state + peak activations + workspace < VRAM
“How long per epoch on MNIST?”	$60K \times \text{forward-backward FLOPs/image/achieved FLOP/s}$
“Is this compute bound or memory bound?”	Compare arithmetic intensity with $R_{\text{peak}}/\text{BW}$

The component-level memory ledger in table 5.8 is explicit but intentionally simplified. In practice, systems engineers work at two levels of fidelity: detailed budgets for design documents and order-of-magnitude estimates for early feasibility gates. The ledger confirms that MNIST fits comfortably in cache while GPT-2 requires a much larger memory tier, before accounting for runtime workspace. A quick mental estimate should reach the same conclusion in seconds, not minutes, and flag any model that cannot physically fit on the target hardware before a single line of profiling code runs.

Feasibility math tells the engineer whether a model *fits*; it says nothing about whether it will *learn*. That depends on how the parameters those bytes hold are first set. Parameter initialization is critical to network behavior. Setting all parameters to zero would cause neurons in a layer to behave identically, preventing diverse feature learning. Instead, weights are typically initialized randomly, often using specific strategies like Xavier/Glorot initialization³² (Glorot and Bengio 2010) or He initialization (He et al. 2015), while biases often start at small constant values or zeros. The scale of these initial values matters: values that are too large or too small lead to poor learning dynamics.

The distribution of parameters affects information flow through layers. In digit recognition, if weights are too small, important input details might not propagate to later layers. If too large, the network might amplify noise. Biases help adjust the activation threshold of each neuron, enabling the network to learn optimal decision boundaries.

Different architectures impose specific constraints on parameter organization. Some share weights across network regions to encode position-invariant pattern recognition; others restrict certain weights to zero, implementing sparse connectivity patterns.

Network architecture, neurons, and parameters are now in place, but a central question remains: the mechanism by which these randomly initialized parameters become useful. A randomly wired network produces outputs no better than chance. Architecture answers *what* the model computes; training answers *how* it learns. The mechanics reveal systems constraints because training adds backward computation, saved activations, gradients, and optimizer state. Batch size then becomes a joint statistical and hardware choice: it changes gradient noise, activation memory, data reuse, and parallel efficiency. The learning process coordinates these costs while transforming 109,386 parameters from random numbers into a functioning digit classifier.

5.4 Learning Process

The MNIST network currently holds 109,386 parameters initialized randomly—numbers that encode no knowledge at all. The transformation of these random values into a digit classifier achieving over 95 percent accuracy relies on four operations repeated millions of times: forward propagation computes a prediction, a loss function measures the error, backpropagation computes each parameter’s loss gradient, and an optimizer uses those gradients to update the parameters.

5.4.1 Supervised learning from labeled examples

A randomly initialized network classifies digits no better than random guessing among ten classes (about 10 percent accuracy). Transforming it into a 95 percent-accurate classifier requires supervised learning: showing the network labeled examples and adjusting its weights based on the errors it makes. Consider the MNIST digit recognition task: the dataset contains 60,000 training images, each a 28×28 pixel grayscale image paired with its correct digit label. The network must learn the relationship between these images and their corresponding digits through an iterative process of prediction and weight adjustment. Ensuring the quality and integrity of training data is essential to model success (chapter 4).

³² | **Xavier/Glorot initialization:** For a layer with n_{in} inputs and n_{out} outputs, a common Glorot choice uses weight variance $2/(n_{\text{in}} + n_{\text{out}})$ to balance forward activations and backward gradients (Glorot and Bengio 2010). The fix costs zero additional FLOPs; it is purely a matter of setting the random distribution at startup.

The relationship between inputs and outputs drives the training methodology. Training operates as a loop where each iteration processes a subset of training examples called a batch.³³ For each batch, the network performs four operations: forward computation through the network layers generates predictions, a loss function evaluates prediction accuracy, weight adjustments are computed based on prediction errors, and network weights are updated to improve future predictions.

The iterative approach can be expressed mathematically. Given an input image x and its true label y , the network computes its prediction according to equation 5.10:

$$\hat{y} = f(x; \theta) \quad (5.10)$$

This equation encapsulates the entire forward pass: the network f takes an input x (say, a 28×28 digit image) and, using its current parameters θ (all the weights and biases we examined earlier), produces a prediction $\hat{y}$ (a vector of ten probabilities, one per digit). The semicolon notation $f(x; \theta)$ distinguishes the input x , which changes with every example, from θ , which remains fixed during inference but evolves during training. The network's error is measured by a loss function³⁴ $\mathcal{L}$ (equation 5.11):

$$\text{loss} = \mathcal{L}(\hat{y}, y) \quad (5.11)$$

The error measurement drives the adjustment of network parameters through backpropagation, examined in section 5.4.4. In practice, training operates on batches of examples rather than individual inputs. For the MNIST dataset, each training iteration might process 32, 64, or 128 images simultaneously for reasons developed in section 5.4.4.6. The training cycle continues until the network achieves sufficient accuracy or reaches a predetermined number of iterations. Throughout this process, the loss function serves as a guide, its minimization indicating improved performance. Establishing proper metrics and evaluation protocols is essential for assessing training effectiveness (chapter 12).

5.4.2 Forward pass computation

An MNIST image becomes ten class scores by moving through weighted layers and nonlinear activations. That computation is **forward propagation**: input data flows through the network's layers to generate predictions. During training, those predictions are then compared with the true value; the resulting loss enters the backward-and-update loop shown in figure 5.17. The forward pass itself underlies both inference and training.




Figure 5.17: Training Loop Architecture: A prediction and true value meet at the loss function; backpropagation computes parameter gradients, and the optimizer turns those gradients into parameter updates for the next forward pass.

Data moves forward through the layers (the red arrow in figure 5.17), while gradients flow backward to update weights (the orange arrow). The figure reveals an important asymmetry: forward propagation produces model outputs, while backward propagation computes gradients for the parameters. *This dependence is why training retains the operation-specific forward values needed by the backward pass.*

³³ | **Batch processing:** Batching serves dual purposes. Larger batches average gradient noise across more examples and can improve hardware utilization by turning matrix-vector work into larger matrix-matrix kernels. They also perform more total work and usually increase latency for an individual request; the throughput gain depends on shape and hardware. Each doubling of batch size roughly doubles activation memory, making batch size a joint optimization and systems decision rather than a purely statistical or hardware choice.

³⁴ | **Loss function:** Formalized by Abraham Wald in statistical decision theory as the “cost” of an incorrect decision, $\mathcal{L}$ quantifies the gap between prediction $\hat{y}$ and ground truth y . The choice of loss function shapes the optimization geometry: it determines the gradient landscape that backpropagation must navigate. A loss with flat regions near incorrect predictions produces weak gradients that stall learning, while a loss with steep gradients near the decision boundary accelerates convergence where it matters most—a systems consequence explored in section 5.4.3.

When an image of a handwritten digit enters the network, it undergoes a series of transformations through the layers. Each transformation combines the weighted inputs with learned patterns to progressively extract relevant features. For the 784-128-64-10 digit classifier, a 28×28 pixel image is processed through multiple layers to ultimately produce probabilities for each possible digit (0–9).

The process begins with the input layer, where each pixel's grayscale value becomes an input feature. For MNIST, this means 784 input values ($28 \times 28 = 784$), each normalized between 0 and 1. These values then propagate forward through the hidden layers, where each neuron combines its inputs according to its learned weights and applies a nonlinear activation function.

Each forward pass through the MNIST network (784-128-64-10) requires substantial matrix operations. The first layer alone performs 100,352 MACs per sample. When processing multiple samples in a batch, these operations multiply accordingly, requiring careful management of memory bandwidth and computational resources. Specialized hardware like GPUs executes these operations efficiently through parallel processing.

5.4.2.1 Individual layer processing

The forward computation through a neural network proceeds systematically, with each layer transforming its inputs into increasingly abstract representations. The digit classifier illustrates this: its transformation process occurs in distinct stages. At each layer, the computation involves two key steps: a linear transformation of inputs followed by a nonlinear activation. The linear transformation applies the same weighted sum operation defined in equation 5.2, but now using notation that tracks layer indices (equation 5.12):

$$\mathbf{Z}^{(\ell)} = \mathbf{A}^{(\ell-1)} \mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)} \quad (5.12)$$

Here, $\mathbf{A}^{(\ell-1)}$ contains the activations from the previous layer (the outputs after applying activation functions), $\mathbf{W}^{(\ell)} \in \mathbb{R}^{n_{\ell-1} \times n_{\ell}}$ is the weight matrix for layer ℓ , and $\mathbf{b}^{(\ell)}$ is the bias vector (broadcast across the batch). The superscript (ℓ) keeps track of which layer each parameter belongs to. This row-vector convention matches the single-sample equation in equation 5.2: each row of $\mathbf{A}$ is one sample, and right-multiplying by $\mathbf{W}$ transforms it to the next layer's width.

Following this linear transformation, each layer applies a nonlinear activation function f (here, f or $f^{(\ell)}$ denotes a generic activation function at layer ℓ ; in section 5.3.2.1, σ referred specifically to the sigmoid function), as expressed in equation 5.13:

$$\mathbf{A}^{(\ell)} = f(\mathbf{Z}^{(\ell)}) \quad (5.13)$$

This process repeats at each layer, creating a chain of transformations:

Input $\rightarrow$ Linear Transform $\rightarrow$ Activation $\rightarrow$ Linear Transform $\rightarrow$ Activation $\rightarrow \dots \rightarrow$ Output

Returning to digit recognition, the pixel values first undergo a transformation by the first hidden layer's weights, converting the 784-dimensional input into an intermediate representation. Each subsequent layer further transforms this representation, ultimately producing a 10-dimensional vector of class scores, one for each possible digit.

5.4.2.2 Matrix multiplication formulation

The complete forward propagation process can be expressed as a composition of functions, each representing a layer's transformation. Formalizing this mathematically builds on the MNIST example.

For a network with N_L layers, we can express the full forward computation as equation 5.14:

$$\mathbf{A}^{(N_L)} = f^{(N_L)} \left(\dots f^{(2)} \left(f^{(1)} (\mathbf{X} \mathbf{W}^{(1)} + \mathbf{b}^{(1)}) \mathbf{W}^{(2)} + \mathbf{b}^{(2)} \right) \dots \mathbf{W}^{(N_L)} + \mathbf{b}^{(N_L)} \right) \quad (5.14)$$

This composition reveals that forward propagation is, at its core, a chain of matrix multiplications interleaved with nonlinear activations. Understanding *why* matrix multiplication dominates AI computation requires examining the arithmetic intensity of each operation.

The nested expression unfolds layer by layer, each step consuming the previous layer's activations as its input:

MatMul

Matrix multiplication exceeds 90 percent of this dense network's forward-pass FLOPs.

1. First layer:

$$\mathbf{Z}^{(1)} = \mathbf{X}\mathbf{W}^{(1)} + \mathbf{b}^{(1)}$$

$$\mathbf{A}^{(1)} = f^{(1)}(\mathbf{Z}^{(1)})$$

2. Hidden layers ($\ell = 2, \dots, N_L - 1$):

$$\mathbf{Z}^{(\ell)} = \mathbf{A}^{(\ell-1)}\mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)}$$

$$\mathbf{A}^{(\ell)} = f^{(\ell)}(\mathbf{Z}^{(\ell)})$$

3. Output layer:

$$\mathbf{Z}^{(N_L)} = \mathbf{A}^{(N_L-1)}\mathbf{W}^{(N_L)} + \mathbf{b}^{(N_L)}$$

$$\mathbf{A}^{(N_L)} = f^{(N_L)}(\mathbf{Z}^{(N_L)})$$

In the MNIST example, the operation dimensions for a batch of B images are as follows:

- Input $\mathbf{X}$: $B \times 784$
- First layer weights $\mathbf{W}^{(1)}$: $784 \times n_1$
- Hidden layer weights $\mathbf{W}^{(\ell)}$: $n_{\ell-1} \times n_\ell$
- Output layer weights $\mathbf{W}^{(N_L)}$: $n_{N_L-1} \times 10$

Arithmetic intensity determines whether a kernel is compute or memory bound. Tiled dense matrix multiplication reuses values, whereas an unfused element-wise operation does little work per byte. Table 5.10 compares their idealized intensities.

Table 5.10: Idealized Arithmetic Intensity of Matrix Multiply vs. Element-Wise ReLU: A tiled dense $N \times N$ matrix multiplication performs $2N^3$ operations while moving at least about $3N^2s$ bytes for s bytes per value, so potential reuse grows with N . A standalone element-wise activation remains at $1/(2s)$ FLOP/byte (0.125 FLOP/byte for FP32); implementations often fuse it with a neighboring kernel to avoid a separate memory round trip.

Operation	Operation Count (O)	Minimum Data Movement	Idealized Intensity	Typical Optimization
Matrix Mul ($N \times N$)	$2N^3$	$3N^2s$ bytes	$\approx 2N/(3s)$	Tiling on CPU/GPU/TPU
Element-wise (ReLU)	N^2	$2N^2s$ bytes	$1/(2s)$	Vectorization and fusion

Systems Perspective 5.4: Why matrix multiplication dominates AI

Modern accelerators have large compute arrays relative to off-chip memory bandwidth. They approach peak arithmetic throughput only when operations expose enough reuse and parallel work. Dense and convolutional layers map well to optimized kernels; custom element-wise sequences are often fused so that their intermediate values remain on chip.

For one sample, the expression $\mathbf{xW}$ is a **General Matrix-Vector Multiply (GEMV)**; batching samples turns the same computation into $\mathbf{XW}$, a **GEMM**. Multidimensional operations like 2D/3D convolutions map onto this same GEMM core through lowering transformations (such as `im2col`), which expand overlapping sliding receptive fields into 2D matrix columns so optimized vendor Basic Linear Algebra Subprograms (BLAS) libraries can execute them at peak hardware efficiency. Matrix operations account for over 90 percent of the floating-point operations in this chapter's dense-network example and dominate many important neural workloads, although the fraction depends on architecture and execution regime. To improve performance, engineers use techniques like blocking and tiling to increase cache or on-chip-memory reuse. This hardware-software co-design principle, designing model architectures around operations that specialized hardware can execute efficiently, is central to modern deep learning. Section C.1.4 provides the detailed treatment of GEMM arithmetic intensity, sparse

matrix formats, and the computational complexity of common layer types needed to optimize these operations in practice.

5.4.2.3 Step-by-step computation sequence

Understanding how these mathematical operations translate into actual computation requires examining the forward propagation process for a batch of MNIST images. This process illustrates how data transforms from raw pixel values to digit predictions.

Consider a batch of 32 images entering the network. Each image starts as a 28×28 grid of pixel values, which flattens into a 784-dimensional vector. For the entire batch, this yields an input matrix $\mathbf{X}$ of size 32×784 , where each row represents one image. The values are typically normalized to lie between 0 and 1.

The transformation at each layer proceeds as a shape-changing sequence; the shapes determine both kernel choice and activation storage. The network first takes the input matrix $\mathbf{X}$ (32×784) and transforms it using the first layer's weights. If the first hidden layer has 128 neurons, $\mathbf{W}^{(1)}$ is a 784×128 matrix, so the computation $\mathbf{XW}^{(1)}$ produces a 32×128 matrix. Each element in this matrix then has its corresponding bias added and passes through an activation function. For example, with a ReLU activation, any negative values become zero while positive values remain unchanged. This nonlinear transformation enables the network to learn complex patterns in the data. The final layer transforms its inputs into a 32×10 matrix, where each row contains 10 class scores. Let z_j denote the raw score (logit) for digit j and let z_k range over all 10 digits. Often, these scores are converted to probabilities using the softmax function in equation 5.6:

$$p(\text{digit } j) = \frac{e^{z_j}}{\sum_{k=1}^{10} e^{z_k}}$$

For each image in the batch, this produces a probability distribution over the possible digits. The digit with the highest probability represents the network's prediction.



Napkin Math 5.2: Counting operations in forward pass

Problem: What is the total arithmetic operation count (O) for one forward pass through the MNIST network ($784 \rightarrow 128 \rightarrow 64 \rightarrow 10$) with batch size 32?

Background: A matrix multiplication of dimensions $(M \times K) \times (K \times N)$ requires $2 \times M \times K \times N$ FLOPs (one multiply and one add per output element, summed over K terms). Bias addition adds $M \times N$ floating-point additions. ReLU activation adds $M \times N$ comparisons, so the table reports total operation-equivalent work rather than pure FLOPs for activation rows.

Solution: Table 5.11 breaks down the operation count layer by layer. The total comes to ~ 7 MOp, or $7 \text{ MOp} \div 32 = \sim 219 \text{ KOp}$ per image.

Systems insight:

1. Layer 1 dominates: The first layer accounts for 91.9 percent of all operations because it processes the largest input (784 dimensions). This is why dimensionality reduction in early layer is so impactful.
2. Compute vs. memory: At batch 1, 219 KOp must amortize a parameter read for only one image, giving an arithmetic intensity of $\sim 0.5 \text{ FLOP/byte}$ under this traffic model. At batch 32, one read of the parameters plus one pass through the batch activations moves at least $\sim 564 \text{ KB}$ and raises the intensity to $\sim 12.4 \text{ FLOP/byte}$. Additional cache misses and intermediate traffic can lower the realized value. The comparison shows why batching improves reuse, but memory- or compute-bound behavior still depends on the machine balance introduced in section D.2.1.
3. Scaling intuition: Doubling the hidden layer widths ($784 \rightarrow 256 \rightarrow 128 \rightarrow 10$) increases the operation count by about $2.1\times$ to about 15 MOp. This comes from recomputing each layer: layers 1 and 3 double, layer 2 quadruples, so the total grows by about $2.15\times$ rather than four times.

Table 5.11: MNIST Forward Pass Operation Count by Layer: Operation-equivalent breakdown for a single forward pass through the 784-128-64-10 network at batch size 32. Matrix multiplication rows are FLOPs; bias, ReLU, and softmax rows are elementwise operations included to show their negligible contribution. Layer 1’s matrix multiply dominates because it multiplies the 784-dimensional input against the widest weight matrix; bias and activation operations contribute under 1 percent of the total.

Layer	Operation	Dimensions	Operations
Layer 1	MatMul	$(32 \times 784) \times (784 \times 128)$	$2 \times 32 \times 784 \times 128 = 6,422,528$
Layer 1	Bias + ReLU	32×128	$2 \times 4,096 = 8,192$
Layer 2	MatMul	$(32 \times 128) \times (128 \times 64)$	$2 \times 32 \times 128 \times 64 = 524,288$
Layer 2	Bias + ReLU	32×64	$2 \times 2,048 = 4,096$
Layer 3	MatMul	$(32 \times 64) \times (64 \times 10)$	$2 \times 32 \times 64 \times 10 = 40,960$
Layer 3	Bias + Softmax	32×10	~ 640 (simplified)
Total			$\sim 7 \text{ MOp}$

5.4.2.4 Implementation and optimization considerations

Forward propagation is easy to state mathematically, but its implementation is constrained by activation storage, batch size, memory layout, and hardware fit. Memory management plays a central role during forward propagation because each layer’s activations must be stored for the backward pass during training. For the MNIST example (784-128-64-10) with a batch size of 32, the activation storage requirements are:

- Input layer: $32 \times 784 = 25,088$ values
- First hidden layer: $32 \times 128 = 4,096$ values
- Second hidden layer: $32 \times 64 = 2,048$ values
- Output layer: $32 \times 10 = 320$ values

This produces a total of 31,552 values that must be maintained in memory for each batch during training, consistent with the worked example in section 5.3.4.3. The memory requirements scale linearly with batch size and become substantial for larger networks.

Batch processing introduces important trade-offs. Larger batches enable more efficient matrix operations and better hardware utilization but require more memory. For example, doubling the batch size to 64 would double the memory requirements for activations. This relationship between batch size, memory usage, and computational efficiency guides the choice of batch size in practice.

The organization of computations also affects performance. Matrix operations can be optimized through careful memory layout and specialized libraries. The choice of activation functions affects both the network’s learning capabilities and computational efficiency, as some functions (like ReLU) require less computation than others (like tanh or sigmoid).

The computational characteristics of neural networks favor parallel processing architectures. While traditional CPUs can execute these operations, GPUs designed for parallel computation can be substantially faster for large dense matrix operations. Specialized AI accelerators achieve even better efficiency through reduced precision arithmetic, specialized memory architectures, and dataflow optimizations tailored for neural network computation patterns.

Energy per inference can vary substantially across hardware platforms. CPUs offer flexible execution, GPUs provide high throughput through parallelism, and specialized edge accelerators trade generality for efficient execution of selected operators. Power alone does not determine energy: the relevant quantity is power integrated over execution time for the same workload and quality target. Memory hierarchy and data movement often explain much of the difference. These considerations recur throughout subsequent chapters, particularly in chapter 6 where architecture-specific optimizations introduce additional trade-offs.

Forward propagation transforms inputs into predictions, but a prediction alone is useless for learning. The training loop requires a way to measure *how wrong* that prediction is in a form that guides weight adjustments. Loss functions fill this role: they translate the gap between prediction and reality into a single number that optimization can minimize.

5.4.3 Loss functions

The forward propagation process described in section 5.4.2 suffices for **inference**, using a pretrained model to make predictions. To train a model, however, we need a way to measure how well those predictions match reality. Loss functions quantify these errors, serving as the feedback mechanism that guides learning. They convert the abstract goal of “making good predictions” into a concrete optimization problem.

Continuing with the MNIST digit recognition example: when the network processes a handwritten digit image, it outputs ten class scores, often converted to predicted probabilities for digits 0–9. The loss function measures those predictions against the target. If an image displays a “seven,” cross-entropy penalizes the model when it assigns too little probability to digit seven. A higher loss indicates a worse value of the chosen objective on that example; it does not by itself diagnose the cause or establish that every task metric is worse.

5.4.3.1 Error measurement fundamentals

A loss function assigns a scalar cost to a prediction and target under a chosen objective. Lower loss indicates a better value of that objective on the evaluated examples, but it is not the same as higher accuracy: confidence and calibration can change while the predicted class remains unchanged. During training, gradients of the loss guide parameter updates. In handwritten-digit recognition, cross-entropy penalizes predictions that assign low probability to the correct digit.

Mathematically, a loss function $\mathcal{L}$ takes two inputs: the network’s predictions $\hat{y}$ and the true values y . For a single training example in digit classification, the loss measures the discrepancy between prediction and truth. When training with batches of data, we typically compute the average loss across all examples in the batch (equation 5.15):

$$\mathcal{L}_{\text{batch}} = \frac{1}{B} \sum_{i=1}^B \mathcal{L}(\hat{y}_i, y_i) \quad (5.15)$$

where B is the batch size and $(\hat{y}_i, y_i)$ represents the prediction and truth for the i -th example. Averaging keeps the loss and mean-gradient scale comparable across batch sizes, although the best learning rate can still change with batch size and optimization regime. The summation also maps naturally to parallel hardware: each example’s loss can be computed independently before a reduction combines them.

The choice of loss function depends on the type of task. For digit classification, the loss function must handle probability distributions over multiple classes, provide meaningful gradients that guide learning, penalize wrong predictions in proportion to their severity, and scale efficiently with batch processing. Cross-entropy loss satisfies all four requirements.

5.4.3.2 Cross-entropy and classification loss functions

For classification tasks like MNIST digit recognition, **cross-entropy loss** is a common way to compare predicted probability distributions with true class labels. The information-theoretic idea of entropy traces to Shannon (Shannon 1948); in supervised classification, cross-entropy penalizes low probability assigned to the correct class.

The logarithm gives **cross-entropy loss** its optimization behavior because probabilities near zero for the correct class incur increasingly large penalties. This shape produces strong gradients for confidently wrong predictions and makes the correct-class probability the objective’s central quantity.

For a single digit image, the network outputs a probability distribution over the 10 possible digits. We represent the true label as a one-hot vector³⁵ where all entries are 0 except for a one at the correct digit’s position. For instance, if the true digit is “seven”, the label would be $y = [0, 0, 0, 0, 0, 0, 1, 0, 0]$.

The cross-entropy loss for this example is defined in equation 5.16:

$$\mathcal{L}(\hat{y}, y) = - \sum_{j=1}^{10} y_j \log(\hat{y}_j) \quad (5.16)$$

where $\hat{y}_j$ represents the network’s predicted probability for digit j . Given our one-hot encoding, this simplifies to equation 5.17:

$$\mathcal{L}(\hat{y}, y) = -\log(\hat{y}_c) \quad (5.17)$$

³⁵ | **One-hot encoding:** Representing K classes as K -dimensional binary vectors where exactly one element is 1; this encoding is sparse by construction: for MNIST’s 10 classes, 90 percent of each label vector is zeros. Implementations commonly store an integer class index instead of materializing that vector. For very large output spaces, methods such as sampled softmax can also reduce the output computation.

where c is the index of the correct class. The loss therefore depends only on the predicted probability for the correct digit; the network is penalized based on how confident it is in the right answer.

For example, if the network predicts the following probabilities for an image of “seven”:

Predicted: [0.1, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.8, 0.0, 0.1]
 True: [0, 0, 0, 0, 0, 0, 0, 1, 0, 0]

The loss would be $-\log(0.8)$, which is approximately 0.223. If the network were more confident and predicted 0.9 for the correct digit, the loss would decrease to approximately 0.105.

5.4.3.3 Batch loss calculation methods

The practical computation of loss involves considerations for both numerical stability and batch processing. When working with batches of data, we compute the average loss across all examples in the batch.

For a batch of B examples, the cross-entropy loss becomes equation 5.18:

$$\mathcal{L}_{\text{batch}} = -\frac{1}{B} \sum_{i=1}^B \sum_{j=1}^{10} y_{ij} \log(\hat{y}_{ij}) \quad (5.18)$$

Here, i indexes examples in the batch, j indexes the ten output classes, y_{ij} is the one-hot target indicator for class j on example i , and $\hat{y}_{ij}$ is the predicted probability for that class. Averaging over B keeps the loss scale comparable as batch size changes.

Computing this loss efficiently requires careful consideration of numerical precision. The dangerous case is not an ordinary small probability, but a probability that an unstable softmax rounds to zero. If the correct-class probability becomes 0, then $\log(0)$ becomes $-\infty$ and subsequent arithmetic can produce a not-a-number value (NaN).

Stable implementations avoid materializing a numerically fragile probability before taking its logarithm:

1. Shift logits before exponentiation, the core of the log-sum-exp trick (section B.2.4), as shown in equation 5.19:

$$\text{softmax}(z_i) = \frac{\exp(z_i - \max(z))}{\sum_j \exp(z_j - \max(z))} \quad (5.19)$$

Subtracting the same $\max(z)$ from every logit leaves the softmax probabilities unchanged because the common factor cancels between numerator and denominator. The numerical effect is decisive: the largest shifted logit becomes zero, so the largest exponential is $\exp(0) = 1$ rather than a potentially overflowing value.

2. Prefer fused cross-entropy from logits. For target class c and $m = \max(z)$, a stable implementation evaluates

$$\mathcal{L} = -(z_c - m) + \log \sum_j \exp(z_j - m) \quad (5.20)$$

Equation 5.20 is mathematically equivalent to $-\log p_c$ but never rounds a tiny softmax probability to zero first. Clamping a probability or adding ϵ can be a defensive fallback, but it changes the objective and should not replace the stable formulation.

5.4.3.4 Impact on learning dynamics

With a batch size of 32 and 10 output classes, each training step processes 32 sets of 10 scores, computes a loss for each example, and averages them into one batch loss. During training, that value serves two roles. Its gradient supplies the update direction, and its trend helps diagnose whether optimization is progressing, stalling, or diverging. Task accuracy, calibration, and other evaluation metrics must still be measured separately on held-out data.

For the MNIST classifier, training loss often begins near $\log(10) \approx 2.3$ when class probabilities are close to uniform, falls rapidly during early optimization, and then improves more slowly. A low training loss indicates confidence on the training examples, not proof of generalization. Validation loss and task metrics determine whether the learned behavior transfers to held-out data.

The loss function's gradients with respect to the network's output logits provide the initial error signal that drives backpropagation. When softmax and cross-entropy are combined, each logit gradient has a particularly simple form: the difference between the predicted and target probabilities, $p - y$. This mathematical property makes cross-entropy loss especially suitable for classification tasks, as it provides strong gradients even when predictions are far from the target.

The choice of loss function also influences other training decisions. Gradient scale interacts with the learning rate, while loss averaging across batches affects gradient noise and thus the useful batch-size regime. Curvature influences optimizer behavior, and convergence decisions require validation metrics and a stopping rule rather than the training-loss trajectory alone.

Loss functions quantify prediction error, but the error signal alone does not tell the system how sensitive that loss is to each parameter. With 109,386 parameters in the MNIST network, estimating every derivative through independent perturbations would be computationally prohibitive. This is the credit assignment problem: connecting the final error to the parameters and intermediate values that can change it. The next section introduces backpropagation, which applies the chain rule to compute all of those local sensitivities efficiently.

5.4.4 Gradient computation and backpropagation

Definition 5.2: Backpropagation

Backpropagation is the gradient-computation algorithm training systems use to apply the chain rule to a computational graph, using recorded operations and saved or recomputed forward-pass values to compute the gradient of the loss with respect to every parameter in a single backward traversal.

1. **Significance:** For dense layers, the backward pass is often on the order of twice the forward-pass FLOPs and requires retaining or recomputing the intermediate values needed by the chain rule. For a model with N_L layers, batch size B , and layer widths $n_1, \dots, n_{N_L}$, a simplified activation ledger scales as $\mathcal{O}(B \sum_{\ell=1}^{N_L} n_\ell)$. Gradients and optimizer state add further memory, so a model that fits on one accelerator for inference may not fit there for training.
2. **Distinction:** Numerical differentiation requires a perturbed evaluation for each of P parameters. Reverse-mode automatic differentiation computes all P gradients in one backward traversal whose cost is proportional to the computational graph, avoiding P separate forward passes.
3. **Common pitfall:** A frequent misconception is that backpropagation is learning. It is a gradient computation algorithm; gradient descent performs the actual parameter update. Confusing the two obscures the systems-level separation: backpropagation determines memory requirements (activation storage), while the optimizer determines additional state requirements (momentum, variance buffers).

A car factory gives a concrete version of the same credit-assignment problem. Vehicles pass through four stations: frame installation (A), engine mounting (B), wheel attachment (C), and final assembly (D). When inspectors find a defective car, they must determine which station caused the problem.

The solution works backward. Starting from the defect, inspectors trace responsibility through each station: how much D's assembly contributed vs. what it received from C, and how much C's work contributed vs. what came from B. Each station receives adjustment feedback proportional to its contribution. If Station B's engine mounting was the primary cause, it receives the strongest signal to change.

Backpropagation follows an analogous backward trace, but it computes sensitivity rather than causal blame. Starting from the loss, it evaluates how a small change in each layer's inputs would change that loss and propagates those derivatives backward. The optimizer later converts the resulting parameter gradients into actual updates.

In neural networks, each layer acts like a station on the assembly line, and backpropagation determines the sensitivity of the final loss to each connection. Translating this intuition into mathematics requires the chain rule of calculus. In the factory analogy, “Station D’s adjustment signal” corresponds to the gradient at the output layer, local sensitivity maps to a partial derivative, and “sending feedback backward” describes the chain-rule multiplication that propagates gradients through the network.

The sensitivity is local to the current parameter values, input batch, and loss. It answers a counterfactual question—how an infinitesimal change here would change this loss—not whether a weight is permanently important or historically caused the prediction. The computational graph makes that question operational. Each node records an operation, and its backward rule maps an incoming loss gradient to gradients for its inputs. The reverse traversal begins at the scalar loss, applies those local rules in reverse dependency order, and accumulates contributions when one value influences the loss through several downstream paths. This accumulation is why backpropagation is more than sending one error signal down a chain: the runtime must respect graph dependencies and combine every relevant path before a gradient is complete. The algorithm below turns that graph-level description into a layer-by-layer procedure.

5.4.4.1 Backpropagation algorithm steps

While forward propagation computes predictions, backward propagation computes the gradients that an optimizer uses to adjust the weights. Consider the running example where the network predicts a “three” for an image of “seven”. Backward propagation systematically calculates how the loss would respond to a small change in each weight at the current parameter values.

The process begins at the network’s output, where the predicted digit probabilities are compared with the true label. The resulting loss gradient then flows backward through the network, giving the local sensitivity of the loss to each layer’s values and weights. The computation follows the chain rule of calculus, breaking down the relationship between weights and loss into manageable steps.

The mathematical foundations of backpropagation provide the theoretical basis for training neural networks, but practical implementation requires software support. Wengert presented automatic numerical derivative evaluation without requiring analytical derivative expressions (Wengert 1964). Modern frameworks expose this capability through automatic differentiation systems that handle gradient computation automatically. Section C.3.1 derives the chain rule formally and shows why **reverse-mode automatic differentiation** computes all parameter gradients in a single backward pass, and section C.3.2 walks through the backward-pass algorithm step by step and explains why it costs roughly twice the FLOPs of the forward pass. Section 7.4.2 later examines the systems engineering aspects of these frameworks. The core implementation contract appears in algorithm 5.1: the forward pass must save the values that the backward pass will need.

Algorithm 5.1 Backpropagation training step

Require: mini-batch inputs $\mathbf{X}$, labels $\mathbf{y}$; N_L layers with parameters $\{\mathbf{W}^{(\ell)}, \mathbf{b}^{(\ell)}\}_{\ell=1}^{N_L}$; loss $\mathcal{L}$

Ensure: parameter gradients $\{\nabla_{\mathbf{W}^{(\ell)}} \mathcal{L}, \nabla_{\mathbf{b}^{(\ell)}} \mathcal{L}\}$ for the optimizer

- 1: $\mathbf{A}^{(0)} \leftarrow \mathbf{X}$
 - 2: **for** $\ell = 1$ to N_L **do** $\triangleright$ *forward pass*
 - 3: compute $\mathbf{Z}^{(\ell)}, \mathbf{A}^{(\ell)}$; save the values its derivative needs
 - 4: **end for**
 - 5: compute loss $\mathcal{L}(\mathbf{A}^{(N_L)}, \mathbf{y})$; seed the output gradient $\nabla_{\mathbf{A}^{(N_L)}} \mathcal{L}$
 - 6: **for** $\ell = N_L$ down to 1 **do** $\triangleright$ *same local rule repeats*
 - 7: use the saved forward values to compute $\nabla_{\mathbf{W}^{(\ell)}} \mathcal{L}, \nabla_{\mathbf{b}^{(\ell)}} \mathcal{L}$
 - 8: propagate $\nabla_{\mathbf{A}^{(\ell-1)}} \mathcal{L}$ to the previous layer
 - 9: **end for**
 - 10: **return** the accumulated parameter gradients
-

Saving activations in the forward loop of algorithm 5.1 is where the systems cost enters: each activation stays live until the backward pass reaches its layer, so **activation memory** scales with batch size, layer count, and activation width. This is why training can exceed inference memory even when the parameter count is unchanged.

5.4.4.2 Error signal propagation

The flow of gradients through a neural network follows a path opposite to forward propagation. Starting from the loss at the output layer, gradients propagate backward, computing how the loss is locally sensitive to each layer and, ultimately, each weight.

Consider what happens when the digit classifier misclassifies a “seven” as a “three”. The loss function generates an initial error signal at the output layer, essentially indicating that the probability for “seven” should increase while the probability for “three” should decrease. This error signal then propagates backward through the network layers.

For a network with N_L layers, the gradient flow can be expressed mathematically. At each layer ℓ , we can express how the layer’s output affected the final loss using the chain rule³⁶ and the schematic Jacobian notation in equation 5.21:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{A}^{(\ell)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{A}^{(\ell+1)}} \frac{\partial \mathbf{A}^{(\ell+1)}}{\partial \mathbf{A}^{(\ell)}} \quad (5.21)$$

This computation cascades backward through the network, with each layer’s gradients depending on those from the layer above it. The result is a local sensitivity, not causal blame: equivalent parameterizations can represent the same function with differently scaled gradients. Backpropagation supplies these sensitivities, while the optimizer may normalize, accumulate, or otherwise transform them before updating the weights.

This process faces challenges in deep networks. As gradients flow backward through many layers, they can vanish or explode. Repeated multiplication can make gradients exponentially small, particularly with saturated sigmoid or tanh activations. Values below the minimum positive *normal* number (1.18×10^{-38} in FP32 or 6.10×10^{-5} in FP16) enter the subnormal range when the format and hardware preserve subnormals; still smaller values underflow to zero, and some accelerator modes flush subnormals to zero for speed. At the other extreme, values beyond the representable maximum can overflow to Inf and propagate NaN, destabilizing training.

5.4.4.3 Derivative calculation process

Computing gradients involves calculating several partial derivatives at each layer: how changes in weights, biases, and activations affect the final loss. These computations follow directly from the chain rule of calculus but must be implemented efficiently for practical training.

At each layer ℓ , we compute three main gradient components. Each serves a distinct purpose in the learning process.

Weight gradients measure the local sensitivity of the final loss to each weight. The optimizer uses these gradients to determine parameter updates (equation 5.22):

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(\ell)}} = \mathbf{A}^{(\ell-1)T} \frac{\partial \mathcal{L}}{\partial \mathbf{Z}^{(\ell)}} \quad (5.22)$$

Bias gradients measure how changing each bias term affects the loss. Since biases shift the activation threshold of neurons, these gradients indicate whether neurons should become more or less easily activated, as expressed in equation 5.23:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{b}^{(\ell)}} = \mathbf{1}^T \frac{\partial \mathcal{L}}{\partial \mathbf{Z}^{(\ell)}} \quad (5.23)$$

Here, $\mathbf{1} \in \mathbb{R}^B$ is an all-ones vector, so left multiplication sums the per-example bias gradients across the batch.

Input gradients propagate the error signal backward to the previous layer. Rather than directly updating parameters, these gradients serve as the “adjustment signals” that allow earlier layers to learn from the final prediction error (equation 5.24):

$$\frac{\partial \mathcal{L}}{\partial \mathbf{A}^{(\ell-1)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{Z}^{(\ell)}} \mathbf{W}^{(\ell)T} \quad (5.24)$$

These three gradient types interact with a practical systems constraint: a desired effective batch may be larger than what fits in accelerator memory at once. Engineers address this through **gradient**

³⁶ | **Chain rule:**
The calculus identity $\frac{\partial \mathcal{L}}{\partial w} = \frac{\partial \mathcal{L}}{\partial a_n} \cdot \frac{\partial a_n}{\partial a_{n-1}} \dots \frac{\partial a_1}{\partial w}$ becomes a *product* of n terms for an n -layer network. If each partial derivative is slightly less than one, the product vanishes exponentially; if slightly greater, it explodes. This multiplicative structure is why depth is a systems constraint, not just a design choice: it dictates the numerical precision requirements and initialization strategies (Glorot and Bengio 2010; He et al. 2015) needed to keep training stable.

accumulation: the full batch is split into smaller *micro-batches* that each fit in memory, and their gradients are accumulated with the appropriate loss normalization before the optimizer step. From the optimizer’s perspective the effective batch size equals the sum across all micro-batches; from the hardware’s perspective each micro-batch is an independent forward-backward pass that stays within the memory budget. For per-example-independent computations, accumulating consistently normalized gradients represents the same batch-gradient sum or mean. Exact numerical equivalence is not universal: floating-point reduction order, batch-dependent operations such as batch normalization, stochastic layers, and batch-dependent losses can make micro-batch execution differ from a single full-batch pass.

Consider the final layer where the network outputs digit probabilities. If the network predicted $[0.1, 0.2, 0.5, \dots, 0.05]$ for an image of “seven”, the gradient flows backward through three steps:

1. Start with the error in these probabilities
2. Compute how weight adjustments would affect this error
3. Propagate these gradients backward to help adjust earlier layer weights

Systems Perspective 5.5: The memory cost of backprop

Activation retention is one component of a broader training-state budget. Under a simplified single-device accounting that excludes temporary workspaces and framework overhead, equation 5.25 captures the four principal byte counts:

$$\text{Training Memory} \approx \text{Model Weights} + \text{Gradients} + \text{Optimizer States} + \text{Activations} \quad (5.25)$$

The derivative equations are also a state-dependency ledger:

- The weight gradient in equation 5.22 needs the incoming activation $\mathbf{A}^{(\ell-1)}$.
- The activation derivative may need $\mathbf{Z}^{(\ell)}$ or a compact mask that records which nonlinear units were active.
- The input gradient in equation 5.24 needs the layer weights $\mathbf{W}^{(\ell)}$.
- The optimizer consumes the completed parameter gradients and may add momentum, variance, or other state after backpropagation has produced them.

A training system can retain these values for the fastest backward pass, recompute selected values to save memory, or offload and partition state across memory tiers or devices. None of these choices eliminates cost: recomputation spends additional FLOPs, while offload and partitioning introduce data movement and synchronization. For deep networks with large batches, wide activations, or high-resolution inputs, this trade-off can dominate accelerator capacity. Section C.3.3 provides the complete training-memory equation and a worked analysis of weights, gradients, optimizer state, and activation costs.

A minimal network makes the gradient arithmetic concrete by tracing actual values through backpropagation.

Napkin Math 5.3: Tracing gradients: A worked backpropagation example

Setup: Consider a network with two inputs, a hidden layer of two neurons (ReLU activation), and one output neuron (no activation, for simplicity). Suppose the current weights and biases are:

- **Hidden layer:** $\mathbf{W}^{(1)} = \begin{bmatrix} 0.5 & -0.3 \\ 0.8 & 0.2 \end{bmatrix}, \mathbf{b}^{(1)} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$

- Output layer: $\mathbf{W}^{(2)} = \begin{bmatrix} 0.6 \\ -0.4 \end{bmatrix}$, $\mathbf{b}^{(2)} = 0$

Given input $\mathbf{x} = [1.0, 0.5]$ and target $y = 1.0$, applying mean squared error: $\mathcal{L} = \frac{1}{2}(\hat{y} - y)^2$.

Forward pass (to establish the values backpropagation needs):

- Hidden preactivation:

$$\mathbf{z}^{(1)} = \mathbf{x}\mathbf{W}^{(1)} + \mathbf{b}^{(1)} = [1.0 \cdot 0.5 + 0.5 \cdot 0.8, 1.0 \cdot (-0.3) + 0.5 \cdot 0.2] = [0.9, -0.2]$$

- Hidden activation (ReLU): $\mathbf{a}^{(1)} = [\max(0, 0.9), \max(0, -0.2)] = [0.9, 0.0]$
- Output: $\hat{y} = \mathbf{a}^{(1)}\mathbf{W}^{(2)} + b^{(2)} = 0.9 \cdot 0.6 + 0.0 \cdot (-0.4) = 0.54$
- Loss: $\mathcal{L} = \frac{1}{2}(0.54 - 1.0)^2 = 0.1058$

Backward pass (applying the chain rule layer by layer):

Step 1: Output layer gradient. The loss gradient with respect to the output is

$$\frac{\partial \mathcal{L}}{\partial \hat{y}} = \hat{y} - y = 0.54 - 1.0 = -0.46.$$

Step 2: Output weight gradients (applying equation 5.22). Since the output layer has no activation function

$$\frac{\partial \mathcal{L}}{\partial z^{(2)}} = \frac{\partial \mathcal{L}}{\partial \hat{y}} = -0.46, \quad \text{and} \quad \frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(2)}} = \mathbf{a}^{(1)T} \frac{\partial \mathcal{L}}{\partial z^{(2)}} = \begin{bmatrix} 0.9 \\ 0.0 \end{bmatrix} \cdot (-0.46) = \begin{bmatrix} -0.414 \\ 0.0 \end{bmatrix}$$

Step 3: Propagate to hidden layer (applying equation 5.24). The error signal sent backward is:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{a}^{(1)}} = \frac{\partial \mathcal{L}}{\partial z^{(2)}} \mathbf{W}^{(2)T} = (-0.46) \cdot [0.6, -0.4] = [-0.276, 0.184]$$

Step 4: Pass through ReLU. The ReLU derivative is one where $z > 0$ and 0 otherwise, so

$$\frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(1)}} = [-0.276 \cdot 1, 0.184 \cdot 0] = [-0.276, 0.0].$$

The second neuron's gradient is zeroed because ReLU blocked its forward signal.

Step 5: Hidden weight gradients (applying equation 5.22):

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(1)}} = \mathbf{x}^T \cdot \frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(1)}} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \cdot [-0.276, 0.0] = \begin{bmatrix} -0.276 & 0.0 \\ -0.138 & 0.0 \end{bmatrix}$$

Weight updates (with learning rate $\eta = 0.1$): Each weight moves opposite to its gradient. For example, $W_{11}^{(1)}$ updates from 0.5 to $0.5 - 0.1 \cdot (-0.276) = 0.5276$, nudging the network toward the correct output. The second hidden neuron's weights receive zero updates for this example because ReLU blocked its activation, illustrating the mechanism that can lead to dead neurons if it happens persistently across the training data.

Systems insight: Backpropagation is not only a calculus procedure; it is also a data-dependency graph. Training systems must preserve the forward activations needed by this backward pass, which is why activation memory becomes a first-order systems cost.

While understanding these mathematical details is essential for debugging and optimization, modern practitioners rarely implement gradients manually. The systems breakthrough lies in how frameworks automatically implement these calculations. Consider a simple operation like matrix multiplication followed by ReLU activation: `output = relu(input @ weight)`. The mathematical

gradient involves computing the derivative of ReLU (0 for negative inputs, 1 for positive) and applying the chain rule for matrix multiplication. The framework records the operation in a computation graph during the forward pass, stores the pre-ReLU activations needed for gradient computation, attaches the backward rule for each operation, and schedules the reverse traversal to balance correctness, memory usage, and hardware utilization. This automation transforms gradient computation from a manual, error-prone process requiring deep mathematical expertise into a reliable system capability that enables rapid experimentation and deployment.

5.4.4.4 Computational implementation details

Activation storage is only the first backward-pass cost. As model size scales, training also adds gradient storage, optimizer-state traffic, and scheduling pressure that the forward-pass memory estimate does not capture.

Consider a larger variant of the MNIST network ($784 \rightarrow 512 \rightarrow 256 \rightarrow 10$) with a batch size of 32. Each layer's activations must be maintained until the backward pass reaches that layer:

- Input layer: 32×784 values (~100 KB using 32-bit numbers)
- Hidden layer 1: 32×512 values (~66 KB)
- Hidden layer 2: 32×256 values (~33 KB)
- Output layer: 32×10 values (~1 KB)

Beyond activations, we must store gradients for each parameter. For this larger network with approximately 535,818 parameters, gradient storage requires a few megabytes. Advanced optimizers like Adam³⁷ roughly triple this by maintaining momentum and velocity terms for every parameter.

Memory bandwidth and footprint compound these capacity requirements. A training step reads parameters, produces gradients, and retains or recomputes the forward-pass values required by each layer's backward calculation. Saved-activation storage generally grows with batch size and network depth, making peak training memory larger than the inference footprint for the same model and batch. For modest networks like the MNIST example, this traffic remains manageable. As models grow, memory bandwidth and activation storage can become binding constraints, motivating high-bandwidth memory, mixed precision, partitioning, and activation checkpointing.

The computational pattern of backward propagation follows a strict sequence: compute gradients at the current layer, update stored gradients, propagate the error signal to the previous layer, and repeat until the input layer is reached. For batch processing, these computations are performed simultaneously across all examples in the batch, enabling efficient use of matrix operations and parallel processing capabilities.

Modern frameworks handle these computations through sophisticated autograd³⁸ engines. Dynamic computation graphs record operations as they execute, while static computation graphs defer execution to expose more optimization opportunities. When a training script asks for gradients, the framework automatically manages memory allocation, operation scheduling, and gradient accumulation across the computation graph. The system tracks which tensors require gradients and schedules operations so that the backward pass follows the dependencies created during the forward pass. This automated management allows practitioners to focus on model design rather than the intricate details of gradient computation implementation.

Checkpoint 5.3: Backpropagation

The credit assignment problem asks how the loss depends on parameters throughout the network. Backpropagation computes those gradients via the chain rule; verify that the mechanism is clear:

Mechanism

- ☐ **Chain rule:** can you explain how equation 5.21 decomposes a complex error signal into local gradients for each layer?
- ☐ **Gradient decay:** why do gradients often vanish in deep networks?

Training vs. inference

³⁷ | **Adam (adaptive moment estimation):** Maintains per-parameter first and second moment estimates, requiring $2 \times$ additional FP32 memory beyond the parameters themselves (Kingma and Ba 2015). For a 100K-parameter MNIST model this overhead is negligible, but for a 7-billion parameter model it adds ~56 GB, often the difference between fitting on one GPU or requiring partitioning. Adam is widely used because its adaptive updates work robustly across many tasks, but learning rate, schedule, regularization, and memory strategy still require validation.

³⁸ | **Autograd (automatic differentiation):** Reverse-mode automatic differentiation traces back to Linnainmaa's work on differentiating computer programs (1970). Modern autograd engines record forward-pass operations into a directed acyclic graph (DAG), then traverse it backward using the chain rule to compute gradients automatically. The key systems trade-off is that more flexible execution captures exactly what happened in each iteration, while more fixed execution exposes more opportunities for ahead-of-time optimization. Chapter 7 develops this framework design choice in detail.

- ❑ **Computational cost:** if the forward pass requires N operations, why does training require roughly $3N$ total operations?
- ❑ **Memory cost:** why must training retain or recompute the forward values needed by the backward pass, and how does this scale with batch size and network depth?
- ❑ **Layer cost:** for a dense layer mapping 4,096 inputs to 4,096 outputs at batch size 32, estimate its forward FLOPs and its FP16 activation bytes to one significant figure, then say which of the two doubling the batch grows.

Backpropagation computes the local gradient of the loss with respect to each weight; the optimizer converts that signal into the actual parameter update. The step size, direction refinement, and momentum across iterations are governed by the optimizer. No optimizer is universally best across all possible problems (Wolpert and Macready 1997), so neural-network training depends on choosing update rules and hyperparameters that match the model, data, and hardware constraints.

5.4.4.5 Parameter update algorithms



Definition 5.3: Gradient descent

Gradient descent is the iterative optimization algorithm that updates model parameters in the direction of the negative gradient, $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$, trading computational steps (O) for loss reduction.

1. **Significance:** Optimizer choice contributes to training memory. In a simplified FP32 ledger, stochastic gradient descent (SGD) needs each weight and gradient (8 bytes per parameter), while Adam adds two moment buffers, bringing those four tensors to 16 bytes per parameter. Mixed-precision training may retain selected FP32 master values while storing other tensors more narrowly; chapter 8 develops the complete accounting. Optimizer state, saved activations, gradients, temporary workspaces, and communication buffers together determine how many accelerators training requires.
2. **Distinction:** Backpropagation computes gradients (the “what to change” signal); an optimizer applies updates (the “change it” action). The backward pass determines which forward values must remain available, while the optimizer adds its own state and update computation.
3. **Common pitfall:** A frequent misconception is that gradient descent finds the global minimum. Neural network loss landscapes are nonconvex with many local minima, saddle points, and plateaus. In practice, SGD and its variants converge to regions of low loss that generalize well, but the path depends on the learning rate schedule, batch size, and initialization.

The optimization process now needs an update rule that turns the backpropagated gradient into parameter changes. This iterative process uses each weight’s local loss sensitivity to update parameters, seeking to reduce loss over successive steps. The fundamental update rule combines backpropagation’s gradient computation with parameter adjustment, as defined in equation 5.26:

$$\theta_{\text{new}} = \theta_{\text{old}} - \eta \nabla_{\theta} \mathcal{L} \quad (5.26)$$

where θ represents the collection of network parameters (weights and biases), η is the learning rate, and $\nabla_{\theta} \mathcal{L}$ is the gradient computed through backpropagation.

For the digit classifier, this means adjusting weights to improve classification accuracy. If the network frequently confuses “seven”s with “one”s, gradient descent will modify weights to better distinguish between these digits. The learning rate η ³⁹ controls adjustment magnitude: too large values cause overshooting optimal parameters, while too small values result in slow convergence.

Despite neural network loss landscapes being highly nonconvex, gradient-based methods often find useful solutions when the model, data, initialization, and optimization regime are well matched. The theoretical reasons, involving concepts like the Lottery Ticket Hypothesis (Frankle

³⁹ | **Learning rate:** This scalar has an outsized impact on training infrastructure because scaling batch size can require retuning the optimization schedule. The linear scaling rule increased learning rate with batch size for a specific large-batch ImageNet training regime (Goyal et al. 2017), but the best policy depends on the model, optimizer, data, and scale. Applying a scaling rule outside its validated regime can cause divergence when moving from single-GPU to multi-GPU training.

and Carbin 2019), implicit bias (Neyshabur et al. 2017), and overparameterization benefits (Nakkiran et al. 2019), remain active research areas. For practical ML systems engineering, the key insight is empirical: learning rates, initialization, regularization, and the training configuration must be validated together.

5.4.4.6 Mini-batch gradient updates

Neural networks typically process multiple examples simultaneously during training, an approach known as mini-batch gradient descent. Rather than updating weights after each individual image, the average gradient over a batch of examples is computed before performing the update.

For a batch of size B , let $\mathcal{L}_i$ denote the loss for example i ; the mean loss gradient becomes equation 5.27:

$$\nabla_{\theta} \mathcal{L}_{\text{batch}} = \frac{1}{B} \sum_{i=1}^B \nabla_{\theta} \mathcal{L}_i \quad (5.27)$$

With a typical batch size of 32, the system performs one forward pass over 32 images, computes 32 per-example losses, averages their gradients, and applies a single weight update. That average smooths noisy per-example gradients, but it also means the accelerator must keep the batch's activations available until the backward pass consumes them. Batch size therefore influences available parallelism and utilization together with model shape, kernel implementation, and hardware.

5.4.4.7 Iterative learning process

The complete training process combines forward propagation, backward propagation, and weight updates into a systematic training loop. This loop repeats until the network achieves satisfactory performance or reaches a predetermined number of iterations.

A single pass through the entire training dataset is called an epoch.⁴⁰ For MNIST, with 60,000 training images and a batch size of 32, each epoch consists of 1,875 batch iterations. Algorithm 5.2 lays out the mini-batch SGD loop: each epoch shuffles the data, steps through it one mini-batch at a time with a forward and backward pass, and applies a single parameter update per batch.

Algorithm 5.2 Mini-batch stochastic gradient descent

Require: training set $\{(x_i, y_i)\}_{i=1}^D$; rate η ; batch size B ; epochs E ; initial θ_0

Ensure: trained parameters θ

```

1:  $\theta \leftarrow \theta_0$ 
2: for  $e = 1$  to  $E$  do
3:   shuffle the training set  $\triangleright$  reduce order-dependent patterns
4:   for each mini-batch  $\mathcal{M}$  of size  $B$  do
5:      $\hat{y}_i \leftarrow f_{\theta}(x_i)$  for all  $(x_i, y_i) \in \mathcal{M}$ 
6:      $\mathcal{L}_{\mathcal{M}} \leftarrow \frac{1}{B} \sum_{(x_i, y_i) \in \mathcal{M}} \mathcal{L}(\hat{y}_i, y_i)$ 
7:      $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}_{\mathcal{M}} \triangleright$  backpropagate, then update
8:   end for
9:   evaluate on validation; stop if the convergence criterion is met
10: end for
11: return  $\theta$ 
```

The mini-batch loop makes the batch size B a hardware unit of work: a larger B can improve matrix utilization and accelerator occupancy but raises activation memory, and because the update fires once per batch rather than once per example, B also sets the gradient-update cadence. The inner loop in algorithm 5.2 is the unit that the hardware repeatedly executes. Forward propagation creates activations for every example in the current mini-batch, backward propagation consumes the required saved values to compute gradients, and the optimizer mutates the parameters once per batch. Training loss tracks the optimization objective on recent batches, while validation accuracy measures task performance on a separate held-out validation set. The test set remains reserved for final evaluation after model and hyperparameter choices are fixed.

⁴⁰ | **Epoch:** One complete pass through a training dataset. The number of passes is a direct multiplier on total work: a 100-epoch MNIST run executes $100 \times$ more forward and backward passes than a single epoch. Large language-model training is often specified by a token budget rather than a simple epoch count because data sources may be mixed and sampled at different rates. The same accounting still applies: every additional training token adds another forward and backward computation.

5.4.4.8 Convergence and stability considerations

A network that achieves 99.5 percent accuracy on training data but only 85 percent on relevant unseen data exhibits a large generalization gap, a strong sign of overfitting. The model may have learned useful patterns while also specializing too closely to the training sample.



Definition 5.4: Overfitting

Overfitting is a generalization failure in which performance on the training sample improves while performance on relevant unseen data does not.

1. **Significance:** The diagnostic signature is measurable: training loss continues to fall while validation loss rises or the relevant validation metric plateaus, creating a widening generalization gap. Parameter count alone does not determine this behavior; data coverage, architecture, optimization, regularization, and the relationship between train and deployment distributions all matter.
2. **Distinction:** Unlike underfitting (where the model lacks the capacity to capture the target function and both training and validation error remain high), overfitting produces low training error but high validation error, indicating that the model has specialized to the training sample rather than learning the underlying distribution.
3. **Common pitfall:** More data helps only when it is representative of the behavior the system must generalize to. Additional examples, augmentation, regularization, and early stopping are possible interventions, but each must be evaluated on held-out data from the relevant operating regime.

Learning rate selection strongly influences training dynamics. A rate that is too large for the optimizer, batch size, normalization, and schedule can make loss oscillate or diverge; one that is too small can waste epochs making negligible progress. There is no universal “moderate” value, so engineers test rates within the intended training configuration and monitor both optimization stability and validation performance.

Convergence monitoring provides essential feedback during training and continues into production deployment, as covered in chapter 14. A flattening training-loss curve signals slowing optimization, but it does not identify the cause: the run may be approaching a useful solution, encountering an optimization limit, or underfitting. A validation plateau is similarly diagnostic rather than conclusive. The gap between training and validation performance helps distinguish overfitting from improvements that transfer to held-out examples. The interplay between batch size, available memory, and computational resources requires careful balancing to achieve efficient training within hardware constraints, the same memory-computation trade-offs established earlier in this section.



Checkpoint 5.4: Neural network learning process

You have now covered the complete training cycle, the mathematical machinery that enables neural networks to learn from data. Before moving to inference and deployment, verify your understanding:

Use the MNIST network ($784 \rightarrow 128 \rightarrow 64 \rightarrow 10$) and a batch size of 32 as the running check.

- ☐ **Forward pass:** can you trace the forward pass using $\mathbf{Z}^{(\ell)} = \mathbf{A}^{(\ell-1)}\mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)}$ and $\mathbf{A}^{(\ell)} = f(\mathbf{Z}^{(\ell)})$?
- ☐ **Cross-entropy loss:** can you explain what cross-entropy loss measures and why batch loss is averaged across examples?
- ☐ **Backward pass:** can you trace how gradients flow backward through the network and why stored activations are needed?
- ☐ **Update rule:** can you write the gradient descent update rule and explain how batch size trades memory, throughput, and gradient stability?

- **Training memory:** can you explain why the full training loop requires more memory than inference?

If any concepts feel unclear, review the earlier sections on forward passes, loss functions, backpropagation, and learning before continuing. These mechanisms form the foundation for the training-vs.-inference distinction that follows.

Selecting a stopping point requires periodically evaluating the model on the validation set and saving its weights to durable storage, a practice called **checkpointing**. For a small model like the MNIST network, writing a checkpoint costs a negligible amount of time. For a model whose parameters occupy tens to hundreds of gigabytes, a synchronous checkpoint write can stall the training loop while those bytes are serialized to disk or object storage; asynchronous and distributed checkpointing reduce that pause but still consume memory, network, and storage bandwidth. Running validation itself requires a full forward pass over the held-out dataset, which at large scale consumes meaningful accelerator time that would otherwise go to training. Checkpointing and validation are therefore not free statistical observations: they impose recurring I/O and compute costs whose frequency must be balanced against recovery and model-selection needs. Saving too infrequently risks losing useful weights after a failure or later degradation; saving too frequently stresses storage bandwidth and adds overhead. This is why convergence monitoring, while conceptually straightforward, becomes a systems infrastructure problem at scale.

The complete training pipeline runs from forward propagation through loss computation to gradient-based weight updates. Training, however, is preparation, not the end goal—once parameters are optimized, deploying the model for real-time predictions changes the resource profile entirely.

5.5 Inference Pipeline

Training transforms randomly initialized weights into parameters that encode useful patterns, but deployment uses those parameters under a different resource profile. During inference,⁴¹ the same mathematical operators face different constraints: latency rather than training throughput, no gradient state, and hardware ranging from edge devices to GPU clusters. Understanding this change is essential for practical systems design.

5.5.1 Production deployment and prediction pipeline

Suppose a model that achieved 99 percent accuracy on its test set begins producing nonsensical outputs three months after deployment, although no code has changed. The weights are frozen and the inference pipeline runs without error. One possible cause is that the operating world moved while the model stood still.

The transition from training to inference limits how a fixed deployed model can adapt. Trained models generalize to unseen inputs through learned statistical patterns, but their parameters do not change during ordinary inference. When operational data diverges from the training distribution, the model continues applying those learned patterns. Consider an autonomous vehicle perception system: if construction zones become more frequent or novel vehicle configurations appear, its responses still reflect the training evidence. Systems that do not update parameters online adapt through monitored data collection, validation, and retraining, a deliberate engineering process detailed in chapter 8.

5.5.1.1 Operational phase differences

Neural network operation divides into two distinct phases with markedly different computational requirements. Figure 5.18 contrasts these phases visually. Inference performs only the forward pass, processing inputs through learned parameters with batch sizes that vary according to demand. Training adds the backward pass for gradient computation and parameter updates, uses batches chosen for optimization and throughput, and must store activations, gradients, and optimizer state simultaneously, consuming significantly more memory.⁴² The learned model may be the same, but execution mode, retained state, and sometimes the deployed graph differ because training-only operations can be removed and inference operations fused or quantized.

⁴¹ | **Inference:** From Latin *inferre* (“to bring in, to conclude”), borrowed from logic where it means deriving conclusions from premises. The ML usage marks a systems boundary: training executes forward and backward passes and maintains optimizer state, while ordinary inference executes the forward graph with frozen parameters. Removing gradients, optimizer state, and most saved activations can reduce memory and computation substantially, but the ratio depends on architecture, batch size, precision, runtime workspace, and serving strategy.

⁴² | **Training GPU power budget:** The “high-memory” requirement is driven by the need to hold parameters, gradients, optimizer state, and activations simultaneously. The corresponding power draw dictates the “substantial cooling infrastructure,” as a single high-end training GPU consumes 400 W–700 W. Even compared with a 4 W mobile inference chip, that is at least 100× the power budget.




Figure 5.18: Inference vs. Training Flow: During inference, neural networks use learned weights for forward pass computation only, simplifying the data flow and reducing computational cost compared to training, which requires both forward and backward passes for weight updates. This streamlined process enables efficient deployment of trained models for real-time predictions.

These computational differences manifest directly in hardware requirements and deployment strategies. Training environments typically employ high-memory accelerators with substantial cooling infrastructure. Inference deployments on constrained hardware prioritize latency and energy efficiency across diverse platforms: mobile devices use low-power neural processors (typically 2–4 W), edge servers use specialized inference accelerators,⁴³ and cloud services often use reduced numerical precision for increased throughput.⁴⁴ Production inference systems serving millions of requests daily require infrastructure concerns, such as request routing and failure handling, that are usually absent from a single training run. At the memory level, training preserves activations for backpropagation, while inference releases layer buffers as soon as possible; table 5.12 turns that difference into a resource profile.

Table 5.12: Training vs. Inference Forward Pass: Both phases execute the same forward equations, but they manage state differently. Training must retain or recompute the values required by the backward pass and also carries gradients and optimizer state. Inference can reuse a rolling pair of layer buffers. At a matched batch size of 32, the simplified MNIST ledger in table 5.8 is about 3.4×; the ratio is workload and implementation dependent.

Characteristic	Training Forward Pass	Inference Forward Pass
Activation Storage	Retains values required by the backward pass	Retains the current input and output buffers
Memory Pattern	Saves or recomputes selected intermediate states	Reuses buffers as successive layers complete
Computational Flow	Structured for gradient computation preparation	Optimized for direct output generation
Resource Profile	Higher memory requirements for training operations	Minimized memory footprint for efficient execution

5.5.1.2 Memory and computational resources

Neural networks consume computational resources differently during inference than during training. Inference has two memory obligations. The first is persistent: the trained weights and biases must remain available for every request. The second is transient: each layer produces an activation buffer that is needed only until the next layer consumes it. This distinction is the reason inference can be much leaner than training, even though the arithmetic in the forward pass is the same.

For the canonical MNIST network (784 → 128 → 64 → 10), the persistent parameter block contains 109,386 parameters, or about 438 KB at 32-bit floating point precision. The layer-level counts in table 5.13 show why: each fully connected layer performs one multiply-add for each weight, so parameter memory and arithmetic scale together.

⁴³ | **Edge inference accelerators:** The Edge TPU (Google Coral) operates in the mobile/embedded tier at about 2 W, delivering 4 TOPS through an INT8 datapath. Edge servers sit in a different tier: Jetson AGX Orin reaches 275 TOPS at 15 W–60 W, about 68.8× more throughput but with wired-power assumptions.

⁴⁴ | **Quantization:** Moving from FP32 to INT8 reduces parameter storage by 4× and can increase throughput on hardware with efficient INT8 datapaths, although the realized speedup depends on kernel support, memory traffic, and workload shape. Inference often permits lower precision than training because it does not need to represent small gradients or optimizer updates across repeated learning steps. The trade-off is not free: aggressive quantization can degrade accuracy, especially on rare or poorly represented inputs, so calibration and task-level evaluation must establish the safe precision for each deployment. Chapter 10 develops quantization techniques in detail.

Table 5.13: MNIST Inference Resources by Layer: The persistent parameter block and multiply-add count for the canonical MNIST network. In fully connected layers, the same weights that occupy memory also determine the arithmetic work per inference.

Layer	Weights	Biases	Multiply-Adds
Layer 1	$784 \times 128 = 100,352$	128	100,352
Layer 2	$128 \times 64 = 8,192$	64	8,192
Output	$64 \times 10 = 640$	10	640
Total	109,386 parameters	Included in total	109,184

This persistent-plus-rolling-buffer model also explains the deployment optimizations that follow. Batching can increase arithmetic reuse and hardware occupancy but expands activation storage. Lower numerical precision can shrink the persistent parameter block and activation buffers,⁴⁵ but only if task quality survives the smaller representation. Hardware-specific layouts can improve cache reuse by keeping working data close to the compute units. Inference makes these optimizations easier to plan because parameters are fixed and many activations have short lifetimes.

5.5.1.3 Performance enhancement techniques

The fixed nature of inference computation presents optimization opportunities unavailable during training. Once parameters are frozen, the predictable computation pattern allows systematic improvements in both memory usage and computational efficiency.

Batch size selection represents a key inference trade-off. During training, large batches stabilized gradient computation, but inference offers more flexibility. Processing single inputs minimizes latency, making it ideal for real-time applications requiring immediate responses. Batch processing, however, improves throughput by using parallel computing capabilities more effectively. For the MNIST network, processing a single image requires storing 202 layer-output activation values (986 values including the input buffer), while a batch of 32 requires 6,464 layer-output activation values but can process more images per unit time on parallel hardware.

Memory management during inference is far more efficient than during training. Since intermediate values serve only forward computation, memory buffers can be reused aggressively. Activation values from each layer need only exist until the next layer's computation completes, enabling in-place operations that reduce the total memory footprint. The fixed nature of inference allows precise memory alignment and access patterns optimized for the underlying hardware architecture.

Hardware-specific optimizations become particularly important during inference. On CPUs, computations can be organized to improve cache utilization and exploit SIMD parallelism. Accelerator deployments benefit from optimized matrix multiplication routines and efficient memory transfer patterns. These optimizations can also reduce energy use and improve hardware utilization, critical factors in real-world deployments.

The predictable nature of inference also enables optimizations like reduced numerical precision. While training typically requires enough floating-point precision to maintain stable learning, inference can often operate with reduced precision while maintaining acceptable accuracy. For the MNIST network, halving the storage width would halve the parameter and activation byte footprint; any speedup would still depend on hardware and kernel support.

These optimization principles, while illustrated through the simple MNIST feedforward network, represent only the foundation of neural network optimization. More sophisticated architectures introduce additional considerations and opportunities, including specialized designs for spatial data, sequences, and context-dependent computation. These architectural variations and their optimizations are explored in chapter 6 and chapter 10. Production deployment considerations, including batching strategies and runtime optimization, are covered in section 13.7 and chapter 14.

5.5.2 Output interpretation and decision making

Neural network outputs become useful only after they are converted back into decisions a conventional system can act on. Preprocessing bridges real-world data into tensor form; postprocessing maps neural outputs into labels, decision thresholds, validation logic, error handling, and downstream messages. In the MNIST running example, logits are not enough: a digit-recognition system needs the most likely digit, a score whose reliability has been evaluated, and a route for uncertain cases to human review or a secondary recognizer.

⁴⁵ | FP32 (single precision):

The IEEE 754 binary32 format uses 32 bits: 1 sign bit, 8 exponent bits, and 23 explicitly stored fraction bits. It remains a common numerical baseline for neural-network computation because its exponent range accommodates values that narrower formats may overflow or underflow. Moving stored tensors from FP32 to FP16 or BF16 halves their byte footprint, while INT8 quarters it. Throughput improves only when the hardware and kernels execute the narrower format efficiently, and integer inference generally requires calibration or quantization-aware training to preserve task quality. See section D.4 for a detailed comparison of numerical formats and their precision-throughput trade-offs.

Those steps have a different performance shape from the forward pass. Inference benefits from batched matrix operations on accelerators, while thresholding, formatting, validation, and exception handling often run as sequential CPU logic. If that surrounding code is ignored, preprocessing and postprocessing can dominate end-to-end latency even when the neural network itself is fast.

The complete neural network lifecycle, from architecture design through training to inference deployment, is a set of mathematical operations with quantifiable resource costs. These operations have so far lived in the controlled environment of our MNIST running example, where data is clean, latency is unconstrained, and hardware is unchallenged. The historical case study tests them against a production deployment, where none of those conditions holds.



Checkpoint 5.5: Complete neural network system

Before examining how these concepts integrate in a real-world deployment, verify your understanding of the complete neural network lifecycle:

Use the MNIST classifier ($784 \rightarrow 128 \rightarrow 64 \rightarrow 10$) as the running production system.

- ☐ **System lifecycle:** can you trace the lifecycle from architecture design to training loop to trained model to inference deployment?
- ☐ **Training vs. inference memory:** at a matched batch size, can you explain why this ledger gives training roughly $3.4\times$ the memory of inference, naming the gradients, optimizer state, and saved-activation terms?
- ☐ **End-to-end trace:** can you trace one digit image through preprocessing, forward propagation, output probabilities, postprocessing, and final prediction?
- ☐ **Bottleneck spotting:** can you identify where bottlenecks might appear in a real-time system processing 100 images per second?
- ☐ **Human-in-the-loop:** can you explain how validated decision thresholds and monitoring determine when human intervention is needed?

Handwritten digit recognition shows how architectural choices, preprocessing, decision thresholds, and human review combine in a working document-processing pipeline.

5.6 Handwritten Digit Recognition

In the late 1980s, LeCun and colleagues demonstrated backpropagation-based recognition of handwritten ZIP-code digits from USPS-supplied images (LeCun, Boser, et al. 1989). Related convolutional document-recognition work later supported commercially deployed bank-check readers (LeCun et al. 1998). These were distinct systems, but they exposed the same engineering pattern: preprocessing normalized variable handwriting, neural inference produced candidate digits, rejection thresholds sent uncertain predictions to human review, and downstream logic converted accepted predictions into operational decisions. This section uses that shared pattern as a systems case study rather than attributing the complete pipeline to one USPS deployment.

5.6.1 The document-recognition challenge

Handwritten ZIP codes provided an early postal research task, while bank checks supplied a documented large-scale deployment setting. Both required models to handle variation in writing style, pen type, stroke thickness, and character formation. Errors carried operational costs, so accuracy alone was insufficient: the system also needed decision scores and validated rejection thresholds that could route uncertain items to human operators. The samples in figure 5.19 illustrate the variability that preprocessing and recognition had to absorb.

The durable systems lesson lies in the pipeline shared by these applications. Image capture and preprocessing constrained the model input, inference produced class scores, rejection thresholds traded automation coverage against error, and human review handled cases outside the accepted confidence range.



Figure 5.19: Handwritten Address Variability: Address and ZIP-code samples exhibit wide variation in stroke width, slant, spacing, and character formation. A document-recognition pipeline must absorb this variation without confusing legitimate writing differences with changes in the intended characters. Samples from the CEDAR CD-ROM 1 USPS Office of Advanced Technology database.

5.6.2 Engineering process and design decisions

Recognizing a clean, centered digit is straightforward. Recognizing handwriting amid variable backgrounds, scale, alignment, and image quality requires engineering decisions at every stage from data preparation to deployment.

The dataset histories must remain separate. Early ZIP-code research used postal digit images supplied for that task (LeCun, Boser, et al. 1989). MNIST was later constructed from National Institute of Standards and Technology (NIST) handwritten-digit databases and standardized into size-normalized benchmark images (LeCun et al. 1998). It did not arise from collecting live envelopes under varying colors, textures, lighting, and orientations. Across both datasets, however, the engineering requirement was the same: preprocessing had to reduce irrelevant variation without erasing the stroke information needed for recognition.

Architecture design balances accuracy, processing time, and memory cost. For the chapter's standardized 28×28 digit input, the central question is how much model capacity the available latency and hardware budget can support.

Training must cover handwriting variation rather than optimize only for a curated test set. Preprocessing normalizes size and orientation, while augmentation can expose the model to additional plausible variation. Evaluation must then measure performance across relevant writing styles and operating conditions, following the systematic workflow described in chapter 3.

Confidence thresholds determine the division of labor between automation and human review. A high threshold rejects more items for manual handling; a low threshold accepts more model errors. An expected-cost rule chooses threshold t to minimize $\text{manual_rate}(t)C_{\text{manual}} + \text{error_rate}(t)C_{\text{error}}$ subject to throughput and latency constraints.

5.6.3 Shared pipeline architecture

Postal ZIP-code research and deployed bank-check readers used different operational systems, but both fit the pattern in figure 5.20. Image acquisition and conventional preprocessing prepare the

input, neural inference produces digit scores, and postprocessing either accepts the result or routes an uncertain item for review. The figure presents this shared architecture as a composite systems pattern.




Figure 5.20: Document-Recognition Pipeline: A composite six-stage flow distilled from postal ZIP-code research and deployed check-reading systems. Image acquisition and conventional preprocessing prepare handwritten digits for neural inference, and raw scores then pass through postprocessing to the final application output.

The pipeline begins with image acquisition. In postal research, the input was an image of a handwritten ZIP-code region; deployed check readers similarly began with scanned document images. Acquisition quality determines how much variation the later preprocessing and recognition stages must absorb.

Preprocessing locates the relevant field, separates or otherwise identifies its characters, and normalizes the resulting images. Thresholding, connected-component analysis, and size normalization are representative conventional techniques. The chapter uses standardized 28×28 images for its MNIST calculations; that dimension is a teaching convention here, not a claim about every postal or check-reading deployment.

The neural network then processes each normalized digit image. The 1989 ZIP-code research system used an early LeNet variant (LeCun, Boser, et al. 1989) with approximately 10,000 parameters, remarkably compact compared with the 109.4K in this chapter’s MNIST network. Both produce digit scores through layered neural computation, although their architectures and operating contexts differ.

Postprocessing converts digit scores into application decisions. For a ZIP-code task, one uncertain digit can send the complete code for further processing or human review. Deployed document readers use the same general principle: accept sufficiently confident predictions and reject uncertain cases instead of forcing an answer.

Deployed document readers must satisfy end-to-end timing constraints spanning acquisition, pre-processing, inference, postprocessing, and downstream action. The exact machinery differs between mail and financial-document applications, but throughput depends on the complete pipeline rather than neural-network latency alone.

5.6.4 Performance outcomes and operational impact

The 1989 ZIP-code study reported its model, recognition, throughput, and training measurements as one experimental system. Table 5.14 keeps those measurements together, revealing how rejection policy and preprocessing affect the meaning of raw classifier speed (LeCun, Boser, et al. 1989).

Table 5.14: Early ZIP-Code Recognition Results: Measurements from the 1989 backpropagation system applied to USPS-supplied digit images. The rejection result shows why an operational error rate must be read together with the fraction of inputs withheld from automatic classification (LeCun, Boser, et al. 1989).

Metric	Reported Result
Error and rejection	1% classification error after rejecting 12.1% of test digits
Throughput	10–12 classifications/sec including acquisition; over 30/sec after normalization
Model parameters	~10,000
Training	23 passes; about 3 days on a Sun-4/260 workstation

The measurements tell a systems story rather than an accuracy story. The ZIP-code experiment showed that learned recognition could handle realistic handwritten digits with an explicit rejection policy. By the late 1990s, related LeNet-based bank-check readers were processing documents at

commercial scale (LeCun et al. 1998). These were separate milestones, joined by the decision to preserve a path for uncertain inputs.

The rejection path, not the raw error rate, is what made the deployment work. Uncertain cases went to human operators rather than forcing every example through the model, so the system bought reliability by choosing where to stop trusting the network.

Performance depends on how closely operating inputs match the variation represented during training. Image quality, writing style, and preprocessing behavior can all shift recognition accuracy, so the physical and statistical pipeline must be evaluated together.

The economic mechanism comes from shifting the operating point, not from eliminating people entirely. Automation handles high-confidence digits, while human operators resolve uncertain cases and preserve a fallback for inputs outside the model's reliable regime.

This history also identifies durable production requirements. Teams must monitor input quality and rejection rates, update training data when operating distributions change, and maintain both the recognition model and its surrounding acquisition and preprocessing machinery. These are lifecycle implications of the combined systems pattern, not documented features of one USPS deployment.

5.6.5 Key engineering lessons and design principles

Across postal ZIP-code research and deployed check readers, the central lesson is that neural computation succeeds only when the surrounding pipeline shares the same operating constraint. Preprocessing, inference, postprocessing, rejection, and downstream action must support the required throughput and error budget together.

This combined history shows why theoretical foundations and systems integration both matter. Converting handwritten marks into reliable operational decisions required coordinated choices in network architecture, training, preprocessing, confidence policy, and downstream handling.

Taken together, the research demonstrations and later document-reader deployments illustrate practices that remain standard: representative training data, confidence metrics, pre- and post-processing, rejection paths, and end-to-end optimization. These operational considerations are formalized in chapter 14.

Modern edge systems operate under different power, latency, and deployment constraints, but the same engineering principles apply: preprocessing for real-world variation, confidence-based routing to human review, and end-to-end pipeline optimization. The combined document-recognition history therefore supplies a reusable systems pattern without treating postal research, check-reader deployment, and MNIST as one application.

Systems Perspective 5.6: Reading historical system comparisons

The early LeNet-class reader and this chapter's MNIST multilayer perceptron show continuity in computation, not a controlled performance benchmark. Their datasets, architectures, software stacks, machines, and operating targets differ. The defensible comparison is structural: both transform normalized pixels into digit scores, both depend on learned parameters and a preprocessing pipeline, and both need an explicit policy for uncertain outputs.

Parameter count supplies one concrete contrast. The 1989 ZIP-code network had approximately 10,000 parameters, while the chapter's fully connected network has 109,386. That difference describes model state; it does not by itself establish latency, energy, throughput, or cost. Those outcomes require a fixed workload and measurements that identify the hardware, numerical precision, runtime, batch size, and end-to-end decision path. Separating these variables prevents a faster deployment from being attributed entirely to silicon when the data, model, software, or operating policy also changed.

This combined history is one instance of a broader alignment pattern: successful ML systems coordinate the task, its data, the algorithm, and the machine. The next section formalizes that pattern.

5.7 D·A·M Taxonomy

Across the document-recognition histories, different resources served different purposes. USPS-supplied ZIP images grounded early recognition research, LeNet supplied an algorithmic mechanism, later bank-check readers supplied a documented deployment setting, and MNIST supplied a standardized benchmark derived from NIST data. They do not form one system, but together they show the three questions in the D·A·M taxonomy: does the data represent the task, can the algorithm learn the required behavior, and can the machine execute it within the operating budget? Appendix A develops the intersections among these axes.

The **Data** axis asks whether the examples and labels describe the behavior the deployed system must recognize. The USPS ZIP-code images, MNIST digits, and bank-check inputs are related but not interchangeable: each embodies a different acquisition process, input distribution, and operating goal. Preprocessing can normalize some variation, but it cannot create missing handwriting styles or repair a label policy that does not match the downstream decision. Data also has a physical footprint. Batch size, image resolution, numerical representation, and storage layout determine how many bytes must enter the training or inference pipeline.

The **Algorithm** axis asks what transformation can learn the desired behavior from that evidence. In the running network, the 784-to-128 first layer determines both the model's initial representational capacity and its dominant parameter count. Activation functions determine which nonlinear boundaries can be expressed and how gradients propagate; the loss and optimizer determine which errors drive parameter updates. These choices shape the work the system must perform, but an elegant learning rule is useful only if its assumptions match the data and its operations fit the deployment budget.

The **Machine** axis asks whether those operations can execute with the required latency, throughput, memory capacity, and energy. Forward propagation moves parameters and activations through matrix operations; backpropagation adds saved values, gradients, optimizer state, and a reverse dependency traversal. Batch size can increase reuse while expanding activation storage. Lower precision can reduce bytes and expose faster kernels, but only after task-level evaluation shows that the resulting numerical error is acceptable. The machine therefore constrains not only how quickly a fixed network runs, but also which network, batch, and precision are feasible.

The pairwise intersections explain why the axes cannot be optimized independently. Data and Algorithm determine *what the system can learn from*; Data and Machine determine *how information is represented and moved*; Algorithm and Machine determine *how the computation executes efficiently*. At the center, all three questions must be answered together. A larger dataset may improve coverage while increasing input traffic, a wider network may improve capacity while exceeding memory, and a lower-precision kernel may improve throughput while changing accuracy.

The chapter's computations make this diagnostic procedure concrete. Begin with an observed symptom such as low accuracy, missed latency, or exhausted memory. Form competing hypotheses on all three axes, then measure the quantity that distinguishes them: coverage and label quality for Data, error slices and learning behavior for Algorithm, or utilization, bandwidth, and capacity for Machine. An intervention on one axis can move the bottleneck to another, so the final test is always end-to-end behavior under the intended workload. Appendix A turns this first-pass diagnosis into a detailed map of the intersection techniques used throughout the book.

These foundations equip engineers to reason about neural networks from first principles. Yet conceptual understanding alone is insufficient: practitioners must also recognize the recurring misconceptions that derail real-world projects.

5.8 Fallacies and Pitfalls

Statistical learning adds failure modes that line-by-line code inspection alone cannot reveal. The following fallacies and pitfalls can lead teams to misallocate effort or deploy inappropriate solutions.

Fallacy: *Neural networks are entirely opaque and cannot be debugged.*

Neural networks do not offer the line-by-line causal trace of traditional code, but they are not beyond diagnosis. Activation visualization can reveal learned patterns, gradient analysis measures local input sensitivity, and ablation studies estimate component contributions. These tools provide partial evidence rather than a complete explanation: saliency can be unstable, correlated features complicate

attribution, and an ablation may change the behavior it is meant to probe. For the MNIST classifier in section 5.3.2, first-layer weights can still provide a useful check on what patterns the model has learned. Debugging a statistical system therefore combines model-level probes with data slices, controlled interventions, and end-to-end evaluation.

Pitfall: *Discarding domain expertise because a deep model is available.*

Teams assume automatic feature learning removes the need for domain knowledge. Successful systems require domain expertise at every stage: architecture selection, training objective design, dataset curation, and output interpretation. The document-recognition history in section 5.6 shows why confidence thresholds must reflect operating economics and route uncertain cases to human operators. Without domain knowledge, teams can deploy networks that look strong on a test set but fail in production because their thresholds, fallback paths, or error costs are wrong.

Fallacy: *Deeper networks are always more accurate than wider ones.*

Engineers assume that stacking more layers is the primary path to higher accuracy, since depth enables hierarchical feature extraction. In practice, depth alone encounters diminishing returns. ResNet showed that very deep residual networks can train effectively, but its ImageNet results also make clear that more layers must be weighed against added training and inference cost (He et al. 2016a). EfficientNet later demonstrated that compound scaling of width, depth, and input resolution can outperform depth-only scaling at a given resource budget (Tan and Le 2019). The lesson is not that depth is bad; it is that teams should profile capacity utilization and scale the architecture dimension that gives the best accuracy per FLOP.

Pitfall: *Using neural networks for problems solvable with simpler methods.*

Teams assume deep learning always performs better. On small datasets or approximately linear relationships, a simpler model such as logistic regression may match or outperform a neural network while requiring less training, memory, and maintenance. The magnitude of that advantage is workload dependent, so teams should benchmark a simple baseline and justify added complexity through measured task-specific gains. Neural networks excel at hierarchical pattern discovery (section 5.2.3) but impose additional overhead. Use them when their representation capacity produces enough benefit to justify the measured systems cost.

Fallacy: *Training data distribution issues can be fixed after model design.*


Teams treat training as mechanically feeding data through architectures. Networks on imbalanced datasets can exhibit poor minority-class performance: a fraud detector with 99:1 imbalance achieves 99 percent accuracy by always predicting “not fraud” while catching zero fraud cases. An unweighted average empirical-risk objective can underweight rare but critical classes unless sampling, weighting, loss design, and evaluation address the imbalance. Teams that skip exploratory data analysis may deploy models with strong aggregate metrics but unacceptable minority-class performance.

Pitfall: *Deploying research models to production without addressing system constraints.*

Production adds constraints absent from research: latency budgets (50–100 ms end-to-end), memory limits (2–4 GB for edge devices), and concurrent loads (100–1,000 requests per second (RPS)). The complete pipeline includes preprocessing, inference, and postprocessing (section 5.5). A model achieving 20 ms inference fails its 50 ms budget when preprocessing adds 25 ms and postprocessing adds 10 ms (55 ms total). Separating model development from system design therefore wastes effort on accuracy while ignoring deployment feasibility.

Fallacy: *More compute automatically means faster training.*

Teams purchase expensive GPUs expecting proportional speedups, but memory traffic or launch overhead may bind performance instead. Arithmetic intensity identifies the binding resource. Under a simple traffic model that reads the parameters once and moves each batch activation once, the MNIST forward pass in table 5.11 reaches approximately 12.4 FLOP/byte at batch 32 and still less at batch 1. This lies far below the ridge points of many high-throughput accelerators. The Roofline Model in section D.2.1 explains why additional peak FLOP/s cannot help until the workload supplies enough reusable work per byte. Caches, fusion, vectorization, and launch cost make measured speedup the final test.



Pitfall: *Extrapolating accuracy improvements without considering diminishing returns.*

Teams may extrapolate a 5-point gain from scaling 10K parameters to 100K into another 5-point gain at 1M. Empirical improvement instead diminishes with scale according to the task, metric, data quality, model family, and scaling regime. Within table 5.5, the ImageNet rows show error falling from AlexNet’s 15.3 percent to ResNet-152’s 3.6 percent while training FLOPs rise from 5×10^{17} to roughly 10^{19} . Teams must therefore measure marginal return in their own regime rather than extrapolate linearly.

These fallacies and pitfalls share a common root: treating statistical behavior as if code behavior alone determined it. Recognizing them early directs attention toward the appropriate measurements and diagnostics.

5.9 Summary

Neural computation translates mathematical formulations into physical machine execution. Linear algebra and calculus determine operation cost under the iron law of ML systems ($T = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}}$). Forward propagation establishes the computational work O , often dominated by matrix multiplications, while backpropagation retains selected intermediate values and adds gradients and optimizer state. The resulting memory volume D_{vol} depends on architecture, batch size, precision, optimizer, and recomputation strategy. These mechanics make memory footprints, arithmetic intensities, and latency limits estimable before deployment.

Neural networks replace some hand-written decision rules and features with parameters learned from data. Weighted sums, nonlinear activations, and gradient-based learning provide recurring operations from which many modern architectures are composed.

Neural network architecture demonstrates hierarchical processing, where layers can transform raw inputs into representations useful for prediction. Training adjusts connection weights through iterative optimization, while inference applies the learned parameters to new data. These phases create different requirements for computation, memory, throughput, and latency. In this chapter’s simplified FP32 ledger, a batch-32 training step requires $\sim 3.4\times$ the memory of batch-32 inference because it adds gradients, optimizer state, and saved activations; the ratio changes with optimizer, precision, checkpointing, and architecture. The document-recognition case study showed how these mathematical principles enter a production pipeline in which preprocessing, neural inference, rejection, and postprocessing must satisfy one latency and reliability budget.



Key Takeaways: The math behind the model

- **Each paradigm shift changes the systems cost of representation:** In this 28×28 digit example, the operation count rises from ~ 100 comparisons (rule-based) through $\sim 8,000$ operations (classical ML) to 109,184 MACs (deep learning)—a $1,091.8\times$ scenario increase that reshapes hardware requirements.
- **Neural networks learn representations from data:** These networks replace some hand-coded features with hierarchical representations learned from examples. That flexibility shifts the burden toward representative data, training compute, and validation.
- **Training and inference often emphasize different objectives:** Training commonly favors efficient throughput across repeated updates, while online inference commonly favors per-request latency; offline inference may favor throughput instead. Batch size links utilization, memory, throughput, latency, and statistical behavior across these regimes.
- **Activations are math and hardware:** ReLU is generally cheaper to implement than sigmoid or tanh, and its unit gradient for positive inputs avoids activation saturation on that branch. Its practical advantage combines implementation cost with optimization behavior.
- **Forward propagation often centers on matrix work:** Dense matrix kernels exceed 90 percent of FLOPs in this chapter’s dense-network example and dominate many important

neural workloads, which is why specialized matrix hardware can substantially outperform general-purpose execution for those operations.

- **Backpropagation needs the path:** Solving credit assignment requires retaining or recomputing the values used by the backward pass, so memory cost often determines whether a model can be trained on a given device and motivates techniques that reduce, recompute, or partition training state.
- **The complete ML pipeline determines end-to-end performance:** Preprocessing, neural computation, and postprocessing all contribute to latency and reliability. The document-recognition examples demonstrated that production success depends on the entire pipeline operating within real-world constraints, not on model accuracy on a test set alone.

The running MNIST example made this escalation tangible: the same 28 by 28 digit that required about 100 comparisons demanded 109,184 MACs in even a modest three-layer network—a 1,091.8× increase that generalizes across the systems dimensions captured in table 5.3. These fundamentals primarily develop the algorithm axis of the D·A·M taxonomy while revealing how algorithmic choices propagate into machine constraints.

The mathematical and systems implications emerge through fully connected architectures. The multilayer perceptrons explored here expose a direct resource law: dense connectivity makes both parameter storage and multiply-accumulate count scale with adjacent layer dimensions. This flexibility comes with computational costs. A 28 by 28 MNIST image contains 784 input values, and the first fully connected layer learns 100,352 weights by connecting every input to 128 neurons. That dense connectivity contains no explicit prior for spatial locality: the model must learn from data which nearby and distant pixel relationships matter, while paying to store and evaluate every connection.


These foundations establish the mathematical and systems vocabulary for reasoning about neural network behavior. The forward-backward propagation cycle, activation function choices, and memory-computation trade-offs recur throughout every subsequent chapter, whether analyzing why certain architectures train faster, why lower-precision approximations preserve accuracy in some layers but not others, or why multi-machine training requires careful coordination. Understanding these fundamentals enables engineers to move beyond treating neural networks as black boxes toward principled system design.

Reading the code reveals what a network is asked to do; reading the math reveals what it will cost to do it. That is why this chapter dwells on the arithmetic. A neural network is a point where the algorithm meets the machine: weights must be stored and moved, dense layers map to matrix operations, and saved activations make claims on memory that can determine whether the model fits. The math is where an algorithm signs its contract with the hardware, committing in advance to operations the machine must execute. Reading that contract lets an engineer estimate likely bottlenecks before profiling confirms where the network is fast and where it stalls.

What's Next: From universal to specialized

Building on the system taxonomy established in chapter 2 and the end-to-end pipeline requirements in chapter 3, this chapter showed how generic fully connected networks turn every input-to-neuron connection into memory volume D_{vol} and compute work O . Real-world data has exploitable structure: images exhibit spatial locality, text contains ordered and contextual dependencies, and time-series data contains temporal patterns. Chapter 6 develops specialized architectures such as convolutions and self-attention that encode different structural assumptions. Those choices can reduce some forms of redundant work or improve reuse, but they also introduce new trade-offs in activation storage, memory access, arithmetic intensity, and accelerator parallelization.

Network Architectures




- 6.1 Architectural Principles
- 6.2 MLPs: Dense Pattern Processing
- 6.3 CNNs: Spatial Pattern Processing
- 6.4 RNNs: Sequential Pattern Processing
- 6.5 Attention: Dynamic Processing
- 6.6 Transformers: Parallel Sequence Processing
- 6.7 Sparse Architectures: RecSys
- 6.8 Shared Building Blocks
- 6.9 Computational Primitives
- 6.10 Architecture Selection Framework
- 6.11 Fallacies and Pitfalls
- 6.12 Summary

Purpose

Why is choosing a neural network architecture an infrastructure commitment as well as a modeling decision?

Selecting a neural network architecture is both a modeling decision and a contract with physics. The commitment begins with the structure in the data. An architecture that matches that structure can learn with fewer parameters and examples, while a poor match spends additional data and machine time compensating for the wrong computational pattern. Convolution, attention, recurrence, and embedding lookup do not merely encode different assumptions about data; they create different patterns of computation, state, and data movement. Those patterns determine whether work parallelizes cleanly, whether intermediate state fits in memory, and whether training and serving remain affordable at the required scale. Two architectures can deliver similar predictive quality yet impose different memory footprints and latency profiles while exposing different opportunities for later efficiency techniques. Benchmark accuracy alone therefore cannot reveal which design remains viable. Architectural commitments also propagate beyond the model. Data pipelines adopt particular input structures, training infrastructure is provisioned around the model's compute profile, serving systems are tuned to its request path, and monitoring inherits its failure modes. Once those dependencies accumulate, replacing the architecture can mean rebuilding much of the surrounding system. An architecture decision is sound only when its predictive gains remain compatible with the full deployment context. In D·A·M terms, architecture is algorithm-machine co-design at its root, with the mathematical graph determining the work the machine must perform and the machine budget limiting which graphs remain practical.





Learning Objectives

- Distinguish computational characteristics of MLPs, CNNs, RNNs, Transformers, and DLRM-style recommenders
- Explain how inductive biases exploit structure in different data types
- Analyze computational complexity and memory scaling across architectural families
- Identify building blocks such as skip connections, normalization, and gating that enable deep training
- Apply the architecture selection framework to match data characteristics with model designs
- Evaluate how compute, memory access, and data movement determine hardware mapping efficiency
- Critique architecture-selection fallacies under latency, bandwidth, and parallelization constraints

6.1 Architectural Principles

A dense model spends parameters and memory on relationships an image rarely needs. Convolutional neural networks (CNNs) encode locality; multilayer perceptrons (MLPs) do not. Matrix multiplication, activation functions, and gradient computation form the “verbs” of neural networks. Architectures assemble those verbs into computational graphs: specialized structures optimized for specific data types and computational constraints. Under the silicon contract (principle 4), every architecture makes an implicit agreement with hardware, trading computational patterns for efficiency on particular problem classes.

Every neural network architecture decides how computation should be organized to match the inherent structure in the data. Images have spatial locality, language has sequential dependencies, and tabular records often lack a fixed spatial or sequential organization. The architecture encodes assumptions about these patterns directly into the computational graph, and those assumptions influence parameter count, hardware utilization, and deployment feasibility. Architecture selection is therefore a systems engineering problem that directly affects the iron law terms: the number of operations O and the volume of data movement D_{vol} . The structural assumptions that each architecture encodes are known as **inductive biases**,¹ and they serve as the unifying concept for this entire chapter.

- CNN
- Transformer
- MLP

Inductive bias from strong (CNN) to weak (MLP): stronger prior, less data needed.

¹ | **Inductive bias:** From Latin *inducere*, “to lead into,” encoding a structural assumption “leads” the model toward a smaller solution space, which is why this concept unifies the entire chapter: every architecture discussed here—multilayer perceptron (MLP), convolutional neural network (CNN), recurrent neural network (RNN), and transformer—is defined by its choice of bias. A CNN’s locality bias cuts parameters by orders of magnitude vs. an equivalent MLP, directly shrinking the iron law’s O and D_{vol} terms, while global attention trades quadratic score computation for long-range connectivity.



Definition 6.1: Inductive bias

Inductive bias is a structural constraint built into a model architecture that restricts the hypothesis space, enabling generalization from finite data by encoding domain-specific assumptions (such as spatial locality or sequential ordering) directly into the computational graph.

1. **Significance:** Inductive bias can reduce the dataset size (D) required for generalization. For one output feature detector, a dense connection uses $\mathcal{O}(N_{\text{pix}} C_{\text{in}})$ parameters, while a shared CNN filter uses $\mathcal{O}(K^2 C_{\text{in}})$ parameters and is applied at every position, costing $\mathcal{O}(N_{\text{pix}} K^2 C_{\text{in}})$ operations. For a 224×224 RGB image, a shared 3×3 filter therefore uses roughly $5,575.1 \times$ fewer parameters than one dense feature detector, reducing parameter storage and often the data required to avoid overfitting.
2. **Distinction:** Unlike regularization (which penalizes hypothesis complexity at training time via L1/L2 terms), inductive bias restricts the hypothesis space at architecture design time: a CNN represents nonlocal functions less directly than local patterns, while regularization discourages complexity through the training objective.
3. **Common pitfall:** A frequent misconception is that stronger inductive bias is always better. A strong locality bias (CNN) excels on spatial data but represents long-range dependencies in language less directly than global attention, which supports long-range dependencies at $\mathcal{O}(S^2)$ score-computation cost.

A CNN encodes an inductive bias of spatial locality: nearby pixels matter more than distant ones. A transformer lets each element attend to any other, enabling long-range relationships at quadratic score-computation cost. These biases help architectures learn efficiently by restricting the space of functions they represent. Without an appropriate bias, a model may require substantially more data and compute to learn the same structure. The unified framework in section 6.10 brings these architectural families together once each bias has appeared in practice.

Machine learning systems face a core engineering trade-off: **representational power vs. computational efficiency**. Under the iron law of ML systems (principle 3), architectural choice is the primary determinant of the operation-count term O . A transformer’s attention mechanism enables global relationships but scales as $\mathcal{O}(S^2)$ operations with sequence length S ; a CNN exploits spatial locality to reduce operations to linear scaling in the number of spatial positions. Matching the right inductive biases to a workload’s data while setting a manageable operation-count budget defines the practice of neural architecture selection.

D Data

A Algorithm

M Machine

Architecture is the algorithm axis of D·A·M: it sets the operation-count budget.

Example 6.1: The ensemble Netflix did not ship (2009)

Context: Netflix awarded its \$1M Grand Prize in 2009 for a 10 percent improvement in recommendation accuracy (Johnston 2012).

Mechanism: Bringing the winning improvement into production required substantial engineering effort just as Netflix’s business was shifting from mailed DVDs to streaming. That shift also produced richer viewing signals and changed what personalization needed to optimize.

Impact: The additional accuracy did not justify the production effort, so Netflix did not deploy the Grand Prize code.

Fix: Netflix retained two algorithms from the earlier Progress Prize and redirected its recommendation work toward streaming-era personalization.

Systems lesson: An offline metric gain is only valuable when its integration cost and objective still match the production system.

Choosing an architectural backbone sets the primary memory layout and compute pattern for downstream hardware. Table 6.1 contrasts the five major neural network families across their spatial/temporal inductive biases, tensor operations, and memory access characteristics.

Table 6.1: Neural Architecture Families: Each family targets a distinct data structure and exposes a distinct system bottleneck. The bottleneck column reads directly from the iron law: MLPs stress memory bandwidth through dense activations, CNNs stress compute throughput because convolution reuses filters across many pixels, RNNs stress sequential dependencies that prevent parallelism, transformers compute quadratic score interactions and may materialize quadratic score storage, and DLRM stresses memory capacity at terabyte scale through embedding tables.

Architecture	Data Type	Core Innovation	System Bottleneck
MLPs	Tabular/Unstructured	Dense connectivity	Memory bandwidth
CNNs	Spatial (images)	Local filters + weight sharing	Compute throughput
RNNs	Sequential (time series)	Recurrent state	Sequential dependencies
Transformers	Relational (language)	Dynamic attention	Quadratic score compute; optional S^2 storage
DLRM	Categorical (recommendations)	Embedding tables	Memory capacity (TB+)

Five specific model architectures recur throughout this book as lighthouse models: consistent reference points that ground abstract concepts in concrete systems reality. These examples are concrete implementations of the Workload Archetypes (Compute Beast, Bandwidth Hog, etc.) introduced in section 2.3.2. To understand *why* these specific models were chosen, consider the history of model evolution through the lens of the Pareto frontier (figure 6.1).

These models serve as canonical workloads for understanding system constraints. Each occupies a distinct position on the trade-off between accuracy and computational cost, as mapped in figure 6.1.




Figure 6.1: The Efficiency Frontier: ImageNet Top-1 Accuracy vs. Computational Cost (GMACs). Representative points are compiled from published model reports for AlexNet, VGG, ResNet, Inception, MobileNetV2, EfficientNet, ViT, Swin Transformer, and ConvNeXt (Krizhevsky et al. 2012; Simonyan and Zisserman 2015; He et al. 2016a; Szegedy et al. 2015; Sandler et al. 2018; Tan and Le 2019; Dosovitskiy et al. 2021; Liu et al. 2021; Liu et al. 2022). The dashed ‘Pareto frontier’ is an illustrative trade-off curve rather than a benchmark rerun; exact accuracy and reported compute values vary by model variant, training recipe, implementation, and operation-counting convention (multiply-accumulates vs. floating-point operations). The legend distinguishes six families: legacy CNNs (gray), dense CNNs (blue), efficient Mobile architectures (green), EfficientNet-class models (violet), modern CNNs (orange), and Transformers (red). The frontier is led by the efficient and modern-CNN families; the Transformers trade substantial compute for flexible long-range modeling without dominating it, and ViT-B/16 sits below the frontier.

The plotted frontier should be read as a historical map of representative architecture papers, not as a controlled benchmark table. It traces the progression from dense CNNs that prioritized accuracy, through MobileNets that reduced compute per unit of accuracy, to transformer architectures that trade substantial computational cost for flexible long-range modeling. Architectural choices made at design time determine where a system lands on this frontier.

6.1.1 Lighthouse roster: Model biographies

Five models earn the lighthouse role because each isolates one system bottleneck that recurs repeatedly: compute (ResNet-50), memory bandwidth (GPT-2), memory capacity (DLRM), edge latency (MobileNetV2), and always-on power for keyword spotting (KWS). The biographies that follow trace each model’s historical context and why it became a useful reference.

ResNet-50 (He et al. 2016a) anchors the compute-intensive vision lighthouse. The Residual Network (ResNet) addressed the degradation problem in very deep plain networks: adding layers could increase training error despite sufficient capacity. By introducing “skip connections” that improve optimization and gradient flow, it enabled networks of 50, 100, or even 1000 layers. The ResNet architecture won the ImageNet 2015 competition (with very deep 152-layer models), and ResNet-50 has become a widely used backbone and benchmark workload for computer vision. From a systems perspective, it is a highly regular, compute-intensive workload composed almost entirely of dense convolutions, making it a useful test for GPU floating-point throughput.

GPT-2 (Radford et al. 2019) anchors the bandwidth-bound language lighthouse. Generative Pre-trained Transformer 2 demonstrated that scaling up a simple decoder-only transformer architecture on massive datasets could produce coherent text generation. Unlike BERT, an encoder-style transformer that reads context in both directions, GPT-2 generates text sequentially (autoregressively²), creating substantial memory bandwidth pressure during batch-one and small-batch decoding because model weights are read for each decoding step. It serves as our archetype for large language models such as Llama and ChatGPT.

DLRM (Naumov et al. 2019) anchors the sparse-recommendation lighthouse. Meta open-sourced DLRM to expose a workload that differs from CNNs and transformers in a critical way. While vision

² **Autoregressive generation:** A decoding strategy where each output token is conditioned on all previously generated tokens, requiring a full model forward pass per token; for a 1.5B-parameter model in FP16, the weight-only batch-one model used here reads about 3 GB of weights per decoding step and yields a work-to-byte ratio of about 1 FLOP/byte. This low arithmetic intensity often makes batch-one and small-batch decoding bandwidth bound; larger batches can amortize weight traffic. During generation, serving systems also retain attention state from earlier tokens; section 6.5 names this state the key-value cache.

and language models are compute-heavy, recommendation systems are memory-heavy. They must look up user and item preferences in massive embedding tables that can reach terabytes in size, creating unique challenges for latency-critical serving (chapter 13). DLRM is a useful benchmark for memory capacity and sparse memory access patterns in the data center.



Lighthouse 6.1: Canonical workloads

In computer architecture, the *microprocessor without interlocked pipelined stages (MIPS)* processor is often used to teach pipelining, not because it is the fastest available chip, but because it is the clearest embodiment of reduced instruction set computer (RISC) principles. Similarly, this book uses ResNet-50, GPT-2, DLRM, MobileNetV2, and KWS as **canonical workloads**. By studying these “lighthouses,” we learn engineering principles about math, memory movement, and locality that remain valid even as the specific “state of the art” model architectures evolve.

MobileNet (Howard et al. 2017) anchors the edge-efficiency lighthouse. MobileNet challenged the trend of ever-larger models by prioritizing efficiency. It popularized depthwise separable convolutions for efficient vision models, an architectural innovation that reduced FLOPs by $8\text{--}9\times$ for 3×3 kernels with minimal accuracy loss. That smaller compute footprint made MobileNet a natural fit for compression and lower-precision deployment techniques [storing and computing with fewer bits per value] covered in chapter 10. It became a reference family for co-designing vision models with the battery and latency constraints of smartphones and embedded devices.

KWS (Warden 2018) anchors the always-on TinyML lighthouse. Keyword spotting models (like those detecting “Hey Siri” or “Ok Google”) represent the extreme end of efficiency. Designed to run on “always-on” microcontrollers with kilobyte-scale memory and milliwatt power budgets, these models (often depthwise separable CNNs) exemplify the constraints of TinyML (Warden and Situnayake 2020; C. R. Banbury et al. 2021; C. Banbury et al. 2021). They force engineers to count every byte and cycle, motivating extreme quantization (INT8 and INT4) and specialized hardware. Together, these biographies establish why the lighthouses are not a model catalog: each one isolates a different bottleneck signature that the arithmetic-intensity analysis can quantify.

6.1.2 Workload signatures: The arithmetic intensity spectrum


ResNet-50 reuses convolutional weights across many spatial positions, while batch-one GPT-2/Llama decode streams large weight and KV-cache state for one token at a time. That contrast motivates arithmetic intensity, the FLOP/byte ratio established in chapter 5 and used with a hardware roofline to assess whether a workload is likely to be compute or memory bound.

These bottlenecks reflect the underlying math, but the calculation here is a weight-only proxy rather than a complete measurement of main-memory traffic. It divides floating-point work by FP32 model-weight bytes, omitting activation, intermediate, and KV-cache traffic. Batching can amortize weight traffic across examples or tokens, while additional state movement can lower realized intensity. Section C.2.1 gives the per-operation FLOP and parameter formulas used in the proxy. Table 6.2 compares the three lighthouse scenarios and exposes a roughly $160.2\times$ gap under these assumptions.

Table 6.2: Lighthouse Workload Signatures: Weight-only FP32 intensity proxies for three model scenarios. These values suggest different compute and bandwidth pressures but do not determine a fixed hardware affinity; realized bottlenecks depend on execution regime and hardware balance.

Model Family	Lighthouse	Intensity (I)	Hardware Affinity
Dense CNN	ResNet-50	~ 80.1 FLOP/byte	Compute-Rich (GPUs/TPUs)
Efficient Vision	MobileNetV2	~ 42.8 FLOP/byte	Balanced (Mobile NPUs)
Transformer	GPT-2 (Inf)	~ 0.50 FLOP/byte	Bandwidth-rich accelerator

This table provides one systems input to architecture selection. Task requirements determine which model families are viable, while batch size, precision, sequence length, implementation,



cache behavior, and target hardware determine whether the proxy predicts the deployed bottleneck. MobileNet often suits constrained vision deployments, and batch-one transformer decoding often places heavy pressure on memory bandwidth, but both choices require measurement in the intended regime.

The “Bottleneck” column in table 6.3 deserves particular attention: it identifies which system resource (compute throughput, memory bandwidth, memory capacity, latency, or power) limits performance for each workload class. In iron law terms (section 1.7), the bottleneck identifies whether O (operations) or D_{vol} (data movement) dominates the runtime. These distinctions determine which optimization strategies prove effective throughout subsequent chapters.

Table 6.3: Lighthouse Model Comparison: Quantitative characteristics and pedagogical roles of the five canonical workloads. Memory is FP32 weight storage, and FLOPs are per inference except for GPT-2 XL, which is quoted per generated token; the DLRM row reports embedding-table entries rather than dense parameters. The bottleneck column indicates the primary system constraint each model reveals: compute-bound models like ResNet stress arithmetic throughput, while bandwidth-bound models like GPT-2 stress memory transfer rates.

Model	Domain	Params	FLOPs/Inf	Memory	Bottleneck	Role in Textbook
ResNet-50	Vision	25.6M	8.2 GFLOP	102.4 MB	Compute	Dense vision throughput
GPT-2 XL	Language	1.5B	3 GFLOP/token	6 GB	Mem. Bandwidth	Token-by-token serving
DLRM	Recommender	25B	Low	100 GB	Mem. Capacity	Embedding tables and capacity planning
MobileNetV2	Edge Vision	3.5M	600 MFLOP	14 MB	Latency	Depthwise convolutions and efficiency
KWS (DS-CNN)	Audio	200K	20 MFLOP	800 KB	Power	Always-on power budget

Architecture selection is ultimately an engineering trade-off between math (O) and memory movement (D_{vol}). The mechanisms behind the signatures in table 6.2 explain why each lighthouse sits where it does on the intensity spectrum. ResNet-50 earns its high intensity because convolutional layers reuse each weight many times across the spatial dimensions of an image (deeper bottleneck layers reach 100–200+ FLOP/byte), so its performance is limited by how fast the hardware can do math. GPT-2 sits at the opposite extreme: each generated token produces only a matrix-vector multiplication rather than the matrix-matrix operations of batch processing, so the system loads massive weights from memory for a single token’s math, and performance is limited by how fast memory can move bits. MobileNet lands between the two at the whole-model level, with individual depthwise layers falling lower once activation traffic is included: depthwise separable convolutions reduce total O but move more data relative to that work, which fits mobile hardware well yet often “starves” high-end GPUs optimized for dense math.



Checkpoint 6.1: Arithmetic intensity and architecture

Match the architectural choice to its systems implication:

- ☐ **Weight reuse (CNNs):** how does reusing the same weights across many inputs change arithmetic intensity?
- ☐ **Large embedding tables (DLRM):** how do massive embedding tables change the balance between data movement and computation?
- ☐ **Autoregressive decoding (GPT-2):** why does token-by-token generation change the batching and bandwidth picture?

This spectrum determines whether the system needs a faster processor or faster memory to improve performance. The framework this book uses to quantify these limits on specific hardware is the **roofline model**, which plots achievable throughput against arithmetic intensity to show where a workload turns from memory bound to compute bound. Section D.2.1 gives the analytical treatment, with applied examples in chapter 11. Section D.2.1.1 works this intensity-to-bottleneck classification

through a real accelerator specification, computing the ridge point of an A100 and showing how the same operation falls on either side of it.

The lighthouse signatures make the next step concrete: inspect each architecture family by the data pattern it targets, the computation it performs, the hardware mapping it induces, and the bottleneck it exposes. This four-part lens ensures that every architecture is evaluated for what it costs to run, not only for what it learns.

6.2 MLPs: Dense Pattern Processing

Consider a smartphone’s spam filter: given a set of features extracted from an email (sender reputation score, number of links, presence of certain keywords), the model must output a single probability: spam or not. This classification task, where every input feature connects to every output, is the domain of fully connected networks. **MLPs**³ represent the fully connected architectures introduced in chapter 5, now examined through the four-part systems lens established earlier.

MLPs embody an inductive bias: *they assume no prior structure in the data, allowing any input to relate to any output*. This architectural choice enables maximum flexibility by treating all input relationships as equally plausible, making MLPs versatile but computationally intensive compared to specialized alternatives. Their computational power was established theoretically by the Universal Approximation Theorem (UAT)⁴ (Cybenko 1989; Hornik et al. 1989), which we encountered as a footnote in chapter 5. This theorem states that a sufficiently large MLP with nonlinear activation functions can approximate any continuous function on a compact domain, given suitable weights and biases. That combination of theoretical universality and dense connectivity is the architectural concept captured by the *multilayer perceptron*.

Definition 6.2: Multilayer perceptrons

Multilayer perceptrons are feed-forward neural network architectures that apply fully connected layers in sequence, where every neuron in one layer connects to every neuron in the next, encoding no structural assumption about the input domain.

1. **Significance:** The lack of structural prior gives dense layers quadratic parameter scaling in layer width: a single layer mapping 1,024 inputs to 1,024 outputs requires 1,048,576 parameters and about 2.1 MB of weight memory in FP16. A 3×3 convolution mapping 1,024 input channels to 1,024 output channels has about 9.4M weights; convolution’s advantage for images comes from spatial weight sharing across positions, not from reducing the channel-mixing matrix itself. This makes MLPs inefficient for high-dimensional structured inputs like images.
2. **Distinction:** Unlike convolutional neural networks, which exploit spatial locality to reduce parameter count, MLPs treat all input elements symmetrically, making them the architecture of choice for tabular data where no spatial or sequential structure is present.
3. **Common pitfall:** A frequent misconception is that MLPs are too simple to matter for complex tasks. MLPs provide a useful dense baseline, but CNNs, recurrent networks, and transformers add operators and connectivity patterns that cannot be reduced to weight sharing alone.

In practice, the UAT explains *why* MLPs succeed across diverse tasks while revealing the gap between theoretical capability and practical implementation. The theorem guarantees that some MLP can approximate any function, yet provides no guidance on requisite network size or weight determination. While MLPs can theoretically solve any pattern recognition problem, doing so may demand impractically large networks or prohibitive computation. This theoretical power drives the selection of MLPs for tabular data, recommendation systems, and problems where input relationships are unknown. At the same time, these practical limitations motivated the development of specialized architectures that exploit data structure for computational efficiency, as section 6.3, section 6.4, and section 6.6 demonstrate.

³ | **Perceptron:** Formed from “perception” and the device suffix “-tron” (as in *cyclotron* and *klystron*), coined by Frank Rosenblatt (Rosenblatt 1957) for the atomic unit of neural computation: a weighted sum followed by a nonlinear activation, extending the earlier McCulloch-Pitts neuron. MLPs are composed entirely of these units arranged in fully-connected layers, so the efficiency of this single operation, a multiply-accumulate, determines system throughput. Modern accelerators execute over 10^{14} of these operations per second, making the perceptron the computational primitive that the entire ML hardware ecosystem is optimized around.

⁴ | **Universal approximation theorem (UAT):** This theorem provides the mathematical guarantee for the MLP’s “no prior structure” inductive bias by proving a sufficiently wide network can approximate any continuous function. The systems-level catch is that “sufficiently wide” can require a number of neurons that grows exponentially with input dimensionality, rendering the theoretical guarantee practically unattainable for even moderately-sized inputs like a 256×256 image.

6.2.1 Learnability gap

The UAT sounds definitive, yet a fundamental gap separates what MLPs can represent from what they can learn in practice. That gap traces to a critical distinction between *what* a network *can represent* and *what* it *can learn*.

Representation capacity refers to the functions an architecture can express given unlimited resources; the UAT established earlier guarantees MLPs have universal representation capacity. This capacity is particularly effective because of the *Manifold Hypothesis*,⁵ which suggests that high-dimensional data actually occupies a much simpler structure. *Learnability* refers to whether gradient descent can find good weights given finite training samples and computational budgets. A function may be representable yet practically unlearnable.

This distinction resolves the apparent paradox of universal approximation and architectural progress. Specialized architectures such as ResNets and transformers improve learnability by embedding inductive biases that match data structure, even when doing so restricts representational capacity.

Three factors create the *learnability gap*:

1. **Sample complexity:** The UAT provides no bounds on training examples needed. For 28×28 images, an MLP treats 784 pixels independently, requiring exponentially many samples to learn spatial correlations. A CNN embeds locality bias, drastically reducing sample requirements. Mathematically, sample complexity can scale exponentially with input dimension for MLPs but polynomially for architectures matching data structure.
2. **Parameter efficiency:** The UAT guarantees some width suffices, but provides no constructive bounds. Some function classes require widths that grow rapidly with input dimension, while architectures with a matching compositional structure can represent them more compactly.
3. **Optimization difficulty:** Even when optimal weights exist, gradient descent may not find them. MLP loss surfaces exhibit complex topology without the regularizing effect of architectural constraints. Specialized architectures reduce the search space, introducing symmetries that gradient descent exploits. The classic MNIST handwritten digit benchmark illustrates this gap between representation and learnability concretely.



Example 6.2: MNIST: Representation vs. learnability

Scenario: Classifying 28×28 MNIST digits using a fully connected MLP versus a CNN.

Diagnosis: A 3-layer MLP requires 20M parameters because it treats all 784 input pixels independently, ignoring spatial structure. A CNN uses local receptive fields and weight sharing, cutting parameters to 421.4K (47× fewer parameters).

Systems lesson: Inductive bias drives parameter efficiency. Incorporating spatial locality into model architecture reduces parameter footprint by orders of magnitude, lowering SRAM memory bandwidth pressure during inference.

⁵ | **Manifold hypothesis:** The assumption that high-dimensional data lies on a low-dimensional surface embedded within the full space; a 256×256 image lives in a 65,536-dimensional space, but “valid cat images” occupy a tiny structured region. Deep networks progressively unfold this crumpled manifold into linearly separable representations. The systems consequence: if data truly occupied the full space, no architecture could learn from feasible dataset sizes; the manifold structure is what makes finite training budgets sufficient.

⁶ | **No free lunch theorem:** Wolpert and Macready’s 1997 result proved that no optimization algorithm outperforms random search across all possible problems: averaged over every conceivable function, all algorithms are equivalent. The ML systems consequence is that an inductive bias (locality, equivariance, attention) can improve performance on problems matching that bias while hurting performance when its assumptions do not hold, making architecture selection an engineering commitment to a problem class.

The learnability gap motivates the core design principle of this chapter: embed inductive biases that match data structure. Each architecture sacrifices theoretical generality for practical learnability. The No Free Lunch theorem⁶ (Wolpert and Macready 1997) formalizes this trade-off: the bias that helps one task may hurt another. CNN’s translation invariance aids image classification but hurts tasks where absolute position matters. Architecture selection is fundamentally the act of matching inductive bias to data structure.

These theoretical insights translate directly into engineering decisions. Appropriate inductive biases reduce parameter counts (enabling edge deployment), accelerate convergence (reducing training costs), and produce structured computation patterns that map efficiently to specialized hardware (chapter 11). A 20M-parameter MLP infeasible for edge deployment becomes a 421.4K-parameter CNN that fits comfortably, a 47× reduction achieved by matching architecture to data structure. The next question is what specific pattern processing requirements dense architectures address.

6.2.2 Pattern processing needs

Deep learning models frequently encounter problems where any input feature may influence any output without inherent constraints. In financial market analysis, any economic indicator may affect any market outcome. In natural language processing, word meaning may depend on any other word in the sentence. These scenarios demand an architectural pattern capable of learning arbitrary relationships across all input features. The architecture must provide unrestricted feature interactions where each output can depend on any combination of inputs, learned feature importance where the system determines which connections matter rather than relying on prescribed relationships, and adaptive representation where the network reshapes internal representations based on the data itself.

The MNIST digit recognition task illustrates this uncertainty concretely. While humans might focus on specific parts of digits (loops in ‘six’ or crossings in ‘eight’), the pixel combinations critical for classification remain indeterminate. A ‘seven’ written with a serif may share pixel patterns with a ‘two’, and variations in handwriting mean discriminative features may appear anywhere in the image. This uncertainty about feature relationships requires a dense processing approach where every pixel can potentially influence the classification decision—an architectural commitment that leads directly to the mathematical foundation of MLPs.

6.2.3 Algorithmic structure

These pattern processing needs demand an architecture capable of relating any input to any output. MLPs solve this with complete connectivity between all nodes. This connectivity requirement manifests through a series of fully-connected layers, where each neuron connects to every neuron in adjacent layers, the “dense” connectivity pattern introduced in chapter 5.

Dense connectivity translates directly into fully connected layers and matrix multiplication operations, the mathematical basis introduced in section 5.4.2.2 that makes MLPs computationally tractable. Figure 6.2 shows how each layer transforms its input through this core operation.




Figure 6.2: Multilayer Perceptron Architecture: A three-layer network in which every neuron connects to all neurons in adjacent layers. The highlighted neuron receives weighted contributions from all inputs, illustrating the dense $\mathcal{O}(N \times M)$ connectivity pattern implemented through matrix multiplications. For MNIST classification, a 784-dimensional input connects to 100 hidden neurons through a 784×100 weight matrix, requiring 78,400 multiply-accumulate operations per sample. Adapted from (Reagan et al. 2017).

Equation 6.1 expresses the dense layer’s affine transformation followed by its activation.

$$\mathbf{h}^{(\ell)} = f(\mathbf{h}^{(\ell-1)}\mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)}) \quad (6.1)$$

$\mathbf{h}^{(\ell)}$ represents the layer ℓ output, $\mathbf{h}^{(\ell-1)}$ represents the input from the previous layer, $\mathbf{W}^{(\ell)}$ denotes the weight matrix for layer ℓ , $\mathbf{b}^{(\ell)}$ denotes the bias vector, and $f(\cdot)$ denotes the activation function; section 5.3.2.2 develops the rectified linear unit (ReLU) and related nonlinearities in detail. This layer-wise transformation, while conceptually simple, creates computational patterns whose efficiency depends critically on how we organize these operations for different problem structures.

The dimensions of these operations reveal the computational scale of dense pattern processing. The input vector $\mathbf{h}^{(0)} \in \mathbb{R}^{d_{\text{in}}}$ (treated as a row vector in this formulation) represents all potential input features. Weight matrices $\mathbf{W}^{(\ell)} \in \mathbb{R}^{d_{\text{in}} \times d_{\text{out}}}$ capture all possible input-output relationships. The

output vector $\mathbf{h}^{(\ell)} \in \mathbb{R}^{d_{\text{out}}}$ produces transformed representations. A four-pixel example turns this bookkeeping into arithmetic.

Example 6.3: Concrete computation example

Consider a simplified 4-pixel image processed by a 3-neuron hidden layer:

Input: $\mathbf{h}^{(0)} = [0.8, 0.2, 0.9, 0.1]$ (4 pixel intensities)

Weight matrix:

$$\mathbf{W}^{(1)} = \begin{bmatrix} 0.5 & 0.1 & -0.2 \\ -0.3 & 0.8 & 0.4 \\ 0.2 & -0.4 & 0.6 \\ 0.7 & 0.3 & -0.1 \end{bmatrix} \quad (4 \times 3 \text{ matrix})$$

Computation:

$$\mathbf{z}^{(1)} = \mathbf{h}^{(0)} \mathbf{W}^{(1)} = \begin{bmatrix} 0.5 \times 0.8 + (-0.3) \times 0.2 + 0.2 \times 0.9 + 0.7 \times 0.1 \\ 0.1 \times 0.8 + 0.8 \times 0.2 + (-0.4) \times 0.9 + 0.3 \times 0.1 \\ (-0.2) \times 0.8 + 0.4 \times 0.2 + 0.6 \times 0.9 + (-0.1) \times 0.1 \end{bmatrix} = \begin{bmatrix} 0.59 \\ -0.09 \\ 0.45 \end{bmatrix}$$

After ReLU: $\mathbf{h}^{(1)} = [0.59, 0, 0.45]$ (negative values zeroed)

Systems insight: Each hidden neuron combines all input pixels with different weights, demonstrating unrestricted feature interaction. Dense layers buy generality by paying the maximum connectivity cost.

The MNIST example makes this scale concrete. The 784-dimensional input connects to every neuron in the first hidden layer. A hidden layer with 100 neurons requires a 784×100 weight matrix (78,400 parameters), where each weight represents a learnable relationship between an input pixel and a hidden feature. This single layer anchors the computational analysis throughout this chapter.

This algorithmic structure enables arbitrary feature relationships while creating specific computational patterns that computer systems must accommodate. Dense connectivity provides the universal approximation capability established earlier but introduces computational redundancy: while the theoretical power of MLPs enables modeling of any continuous function given sufficient width, this flexibility requires numerous parameters to learn relatively simple patterns. Every input feature influences every output, yielding maximum expressiveness at the cost of maximum computational expense. These trade-offs motivate later compression strategies that reduce computational demands while preserving model capability, and chapter 11 explores hardware-specific implementations that exploit regular matrix operation structure.

6.2.4 Computational mapping

Section 6.2.3 defines *what* an MLP computes; computational mapping reveals *how* that computation translates to hardware operations. Listing 6.1 demonstrates how this mapping progresses from mathematical abstraction to computational reality.

Listing 6.1: MLP Layer in Matrix Form: Framework-level matrix operations hide $\mathcal{O}(B \times d_{\text{in}} \times d_{\text{out}})$ multiply-accumulate complexity behind a single function call, enabling hardware-optimized BLAS libraries to use dense matrix hardware efficiently.

```
def mlp_layer_matrix(X, W, b):
    """MLP forward pass using framework-level matrix operations."""
    # X: input matrix (batch_size by num_inputs)
    # W: weight matrix (num_inputs by num_outputs)
    # b: bias vector (num_outputs)

    # Single GEMM call: frameworks dispatch to optimized BLAS/cuBLAS
    # For MNIST: 784 * 100 = 78,400 MACs per sample
    H = activation(matmul(X, W) + b)
    return H
```

The function `mlp_layer_matrix` directly mirrors the mathematical equation, employing high-level matrix operations (`matmul`) to express the computation in a single line while abstracting the underlying complexity. This implementation style characterizes deep learning frameworks, where optimized libraries manage the actual computation.

To understand the system implications of this architecture, we must look “under the hood” of the high-level framework call. The elegant one-line matrix multiplication `output = matmul(X, W)` is, from the hardware’s perspective, a series of nested loops that expose the true computational demands on the system. This translation from logical model to physical execution reveals critical patterns that determine memory access, parallelization strategies, and hardware utilization.

The second implementation in listing 6.2 exposes the actual computational pattern through nested loops, revealing what really happens when we compute a layer’s output: we process each sample in the batch, computing each output neuron by accumulating weighted contributions from all inputs. This translation from mathematical abstraction to concrete computation exposes how dense matrix multiplication decomposes into nested loops of simpler operations. The outer loop processes each sample in the batch, while the middle loop computes values for each output neuron. Within the innermost loop, the system performs repeated multiply-accumulate operations,⁷ combining each input with its corresponding weight.

Listing 6.2: Dense Layer Computation: Nested loops reveal $\mathcal{O}(B \times d_{\text{out}} \times d_{\text{in}})$ complexity where each output neuron requires `num_inputs` multiply-accumulate operations, so the reference MNIST input layer requires 78,400 MACs per sample.

```
def mlp_layer_compute(X, W, b):
    """Explicit loop structure exposing MLP computational patterns."""
    # Loop 1: Process each sample independently (parallelizable)
    for batch in range(batch_size):
        # Loop 2: Compute each output neuron
        for out in range(num_outputs):
            Z[batch, out] = b[out] # Initialize with bias

            # Loop 3: Accumulate weighted inputs (innermost loop)
            # This is the MAC operation: result += input * weight
            for in_ in range(num_inputs):
                Z[batch, out] += X[batch, in_] * W[in_, out]
            # Total per output: num_inputs MACs +
            # num_inputs memory reads

    H = activation(Z) # Element-wise nonlinearity
    return H
```

In our reference MNIST layer, each output neuron requires 78,400 MACs divided by 100, or 784, multiply-accumulate operations and at least 1,568 memory accesses (784 for inputs, 784 for weights). Production implementations call optimized matrix libraries such as Basic Linear Algebra Subprograms (BLAS),⁸ but the same nested-loop pattern still determines the system design problem. The hardware architectures that accelerate these matrix operations, including GPU Tensor Cores⁹ and specialized AI accelerators, are covered in chapter 11.

6.2.5 System implications

Section 6.2.4 showed how MLP operations decompose into nested loops of multiply-accumulate operations. The system-level constraints that emerge from these patterns span three dimensions: memory requirements, computation needs, and data movement.

For dense pattern processing, the memory, compute, and data-movement costs all come from the same source: all-to-all connectivity. Memory usage is dominated by parameter storage. Our reference MNIST layer (784×100) requires only 78,400 parameters, but this $\mathcal{O}(M \times N)$ scaling becomes prohibitive for high-dimensional inputs. A typical 2048-unit layer connected to a 2048-unit layer requires 4,194,304 parameters (16.8 MB at FP32). Since every weight is used exactly once per input vector, there is no opportunity for weight reuse within a single sample processing, making the workload heavily dependent on memory capacity and bandwidth.

7 | MAC (multiply-accumulate): The atomic operation of neural networks: multiply two values and add to a running sum. In Horowitz’s 45 nm reference, an FP32 multiply plus add costs about 4.6 pJ, while a 32-bit off-chip DRAM access costs about 640 pJ (Horowitz 2014). These technology-specific values illustrate why data movement can dominate arithmetic energy; they are not universal constants for current hardware.

8 | BLAS (basic linear algebra subprograms): This standard API for matrix operations enables the use of highly optimized libraries (for example, cuBLAS) to accelerate the 784 multiply-accumulates per neuron. The 784×100 matrix in the MNIST example may use the hardware less efficiently than larger, well-aligned transformer matrices because utilization depends on matrix shape, batching, datatype, library, and accelerator.

9 | Tensor Cores: Specialized units in NVIDIA GPUs that accelerate the thousands of multiply-accumulate operations described by fusing them into single, highly parallelized matrix instructions. Tensor Cores are most efficient when matrix dimensions meet datatype- and architecture-specific alignment multiples; modern cuBLAS/cuDNN can still use Tensor Cores for many nonaligned cases, often with lower efficiency or internal padding. The architectural lesson is vendor-independent: specialized matrix units reward dense, aligned general matrix multiplication (GEMM) workloads and penalize small, irregular shapes that cannot keep the units full.

The core computation is dense GEMV, or GEMM when batched. This computation is regular and parallelizable, but the arithmetic intensity (FLOP/byte) is low for small batches. The **batch size** is the number of input samples processed together; larger batches amortize weight-loading cost over more computations. Modern processors optimize dense layers with single instruction, multiple data (SIMD) units such as AVX-512 on CPUs or systolic arrays on Tensor Processing Units (TPUs) and GPUs, amortizing control overhead over large blocks of parallel arithmetic.

The resulting bottleneck is data movement. To compute 100 hidden values from 784 inputs, the system must move 784×100 weights from memory to the compute units. Applying the arithmetic intensity framework from section 6.1.2 to this layer yields roughly 0.5 FLOP/byte for batch-one FP32 execution. Under a weight-only model that counts one FP32 read per weight and omits input, output, bias, and cache traffic, equation 6.2 gives this ratio, where M and N are the input and output widths.

$$\text{Intensity} \approx \frac{2 \cdot M \cdot N \text{ FLOPs}}{4 \cdot M \cdot N \text{ bytes}} = 0.5 \text{ FLOP/byte} \quad (6.2)$$

On accelerators with ridge points in the hundreds of FLOP/byte, this batch-one FP32 layer is memory-bandwidth bound. Batching can raise intensity by amortizing weight traffic, but the crossover depends on matrix shape, precision, implementation, and hardware. This is why fully connected layers can become inference bottlenecks despite performing fewer total FLOPs than convolutional layers.

Dense connectivity thus moves maximum data for minimum compute. For data with inherent structure (such as spatial locality in images or temporal order in sequences), specialized architectures can exploit that structure for both better accuracy and better efficiency. The most established such architecture is the convolutional neural network.

6.3 CNNs: Spatial Pattern Processing

The MLP's assumption that every input may interact directly with every output proves particularly costly for spatially structured data like images. As the earlier MNIST comparison demonstrated, the example CNN uses $47\times$ fewer parameters by exploiting spatial locality rather than treating every pixel independently.

Convolutional neural networks (CNNs)¹⁰ emerged as the solution to this challenge (LeCun et al. 1998; Krizhevsky et al. 2012). Consider what happens when viewing a photograph: the visual system does not perceive every pixel simultaneously in relation to every other pixel. Instead, it detects local patterns (edges, textures, corners) and composes them into objects. CNNs encode this same insight architecturally.

Spatial locality produces two key innovations that enhance efficiency for spatially structured data. Parameter sharing allows the same feature detector to be applied across different spatial positions, reducing parameters from millions to thousands while improving generalization. Local connectivity restricts connections to spatially adjacent regions, reflecting the insight that spatial proximity correlates with feature relevance. Together, these innovations define convolutional neural networks as an architectural family.



Definition 6.3: Convolutional neural networks

Convolutional neural networks (CNNs) are architectures that exploit translation equivariance and spatial locality to share learned filters across all spatial positions, decoupling parameter count from input resolution.

1. **Significance:** Weight sharing produces dramatic parameter reduction. A 3×3 convolutional layer with 64 input and 64 output channels requires $3 \times 3 \times 64 \times 64 \approx 37,000$ parameters regardless of whether the input image is 224×224 or 1024×1024 . An equivalent fully connected layer on a $224 \times 224 \times 64$ input would require $224^2 \times 64 \times 64 \approx 205$ million parameters, a roughly $5,575\times$ difference. This constant-parameter scaling enables CNNs to process high-resolution inputs within the memory budget of a single accelerator.
2. **Distinction:** Unlike MLPs, which connect every input element to every output element (global connectivity), CNNs restrict each output to a local spatial neighborhood, encoding



Conv 37K

log scale

Weight sharing spares convolution the fully-connected parameter explosion.

¹⁰ | **Convolution:** From Latin *convolvere* (“to roll together”), describing a filter that slides across an input, combining local elements at each position. This “rolling together” enforces a locality constraint that is the source of the operation’s efficiency: a single 5×5 kernel reuses its 25 weights at every spatial position, reducing one feature detector for a 1-megapixel single-channel image from roughly 1,000,000 weights to 25, about $40,000\times$ fewer parameters than a fully connected detector.

the assumption that nearby pixels are more relevant than distant ones. This restriction eliminates entire hypothesis classes at architecture design time rather than penalizing them during training.

3. **Common pitfall:** A frequent misconception is that CNNs are vision-only models. The convolution operation applies to any data with a grid-like topology: 1D convolutions process audio waveforms and time series, 2D convolutions process images and spectrograms, and 3D convolutions process video and volumetric data.

The trade-off is explicit: CNNs sacrifice the theoretical generality of MLPs for practical efficiency gains when data exhibits known structure. Where MLPs treat each input element independently, CNNs exploit spatial relationships to achieve both computational savings and improved accuracy on vision tasks.

6.3.1 Pattern processing needs

Spatial pattern processing addresses scenarios where the relationship between data points depends on their relative positions or proximity. Consider processing a natural image: a pixel's relationship with its neighbors is important for detecting edges, textures, and shapes. These local patterns then combine hierarchically to form more complex features: edges form shapes, shapes form objects, and objects form scenes. The pipeline in figure 6.3 gives this hierarchy a concrete visual form.




Figure 6.3: Spatial Feature Extraction: A zebra image passes through repeated convolution, ReLU, and pooling stages that form progressively deeper feature maps, followed by flattening, fully connected layers, and a softmax classification over Horse, Zebra, and Dog.

This hierarchical processing appears across many domains: local pixel patterns forming edges that combine into objects (computer vision), nearby time-segment correlations identifying phonemes (speech), proximate sensor correlations (sensor networks), and tissue pattern recognition (medical imaging). The approach succeeds not because it mimics the brain, but because it mirrors the compositional structure of the data itself.

Focusing on image processing to illustrate these principles, if we want to detect a cat in an image, certain spatial patterns must be recognized: the triangular shape of ears, the round contours of the face, the texture of fur. These patterns maintain their meaning regardless of where they appear in the image. A cat is still a cat whether it appears in the top-left or bottom-right corner. This indicates two key requirements for spatial pattern processing: the ability to detect local patterns and the ability to recognize these patterns regardless of their position.¹¹ As figure 6.3 illustrates, convolutional neural networks meet both requirements through hierarchical feature extraction, where simple patterns compose into increasingly complex representations at successive layers. CNNs put these spatial processing principles into practice through parameter sharing, local connectivity, and translation equivariance,¹² the key innovations pioneered by Yann LeCun¹³ and LeCun, Boser, et al. (1989).

¹¹ | **ImageNet:** The dataset that validated these two spatial processing requirements at scale. AlexNet's 2012 victory in the ImageNet Large Scale Visual Recognition Challenge reduced top-5 error from 26.2 percent to 15.3 percent on the 1000-class challenge with roughly 1.2 million training images (Krizhevsky et al. 2012); the original ImageNet release contained 3.2 million images across 5,247 synsets (Deng et al. 2009). Subsequent architectures improved accuracy through changes in architecture, optimization, data, and compute budgets, illustrating the interaction between inductive bias and infrastructure cost.

¹² | **Translation equivariance:** An inherent property of the convolution operation where shifting the input guarantees a corresponding spatial shift in the resulting feature map. This is distinct from *invariance*, which pooling or later aggregation can approximate by discarding precise positional data. For example, 2×2 pooling with stride 2 reduces the number of spatial elements by 75 percent.

¹³ | **Yann LeCun and LeNet:** LeCun's architecture directly addressed the intractable scaling of applying dense networks to images by enforcing the principles of local connectivity and parameter sharing. These constraints reduced the parameter count for an image-like input layer by over 95 percent, enabling LeNet-5 to achieve production-grade accuracy on commercial tasks like check reading with only ~60,000 total parameters.

6.3.2 Algorithmic structure

Equation 6.3 sums the local, channel-wise filter products at each output position.

$$\mathbf{H}_{i,j,k}^{(\ell)} = f \left(\sum_m \sum_n \sum_c \mathbf{W}_{m,n,c,k}^{(\ell)} \mathbf{H}_{i+m,j+n,c}^{(\ell-1)} + \mathbf{b}_k^{(\ell)} \right) \quad (6.3)$$

This equation describes how CNNs process spatial data. $\mathbf{H}_{i,j,k}^{(\ell)}$ is the output at spatial position (i,j) in channel k of layer ℓ . The triple sum iterates over the filter dimensions: (m,n) scans the spatial filter size, and c covers input channels. $\mathbf{W}_{m,n,c,k}^{(\ell)}$ represents the filter weights, capturing local spatial patterns. Unlike MLPs that connect all inputs to outputs, CNNs only connect local spatial neighborhoods.

Breaking down the notation further, (i,j) corresponds to spatial positions, k indexes output channels, c indexes input channels, and (m,n) spans the local receptive field.¹⁴ Unlike the dense matrix multiplication of MLPs, this operation applies the same filter weights at each spatial position.

Convolutional layers process local neighborhoods (typically 3×3 or 5×5), reuse the same weights at each spatial position, and maintain spatial structure in the output. The sliding-window mechanics in figure 6.4 show a small filter moving across the input image and computing a dot product at each position to generate a feature map. This operation captures local structures while maintaining translation equivariance—the same filter detects the same pattern regardless of where it appears. For an interactive visual exploration of convolutional networks, the CNN Explainer (Wang et al. 2021) project provides an insightful demonstration of how these networks are constructed.

¹⁴ | **Receptive field:** The input region influencing a particular output neuron. With 3×3 filters, receptive fields grow by 2 pixels per layer, so a neuron at layer 3 “sees” a 7×7 region. This growth rate constrains architecture depth: detecting objects spanning 100+ pixels in a 224×224 image requires either deep stacks of small filters (more layers, more memory for activations) or larger kernels (more parameters per layer), a fundamental depth-vs.-width trade-off in CNN design.




Figure 6.4: Convolution Operation: A 3×3 orange receptive field in a dashed 7×7 input maps to one cell of the violet 5×5 valid-convolution output. Applying a 3×3 filter to that receptive field requires 9 multiply-accumulate operations per input channel for each output value. For a 224×224 RGB input, parameter sharing reduces a dense 150,528-weight input connection to a 27-weight local filter, or roughly $5,575 \times$ fewer weights, while encoding translation equivariance: the same edge detector works at any image location.

To illustrate, consider applying a CNN to the same MNIST images used in our MLP analysis. Each convolutional layer applies a set of filters (for example, 3×3) that slide across the 28×28 input, computing local weighted sums. With 32 filters and padding to preserve dimensions, the layer produces a $28 \times 28 \times 32$ output, where each spatial position contains 32 different feature measurements of its local neighborhood. This contrasts sharply with the MLP approach, where the entire image is flattened into a single vector before processing.

This algorithmic structure directly implements the requirements for spatial pattern processing, creating distinct computational patterns that influence system design. Unlike MLPs, convolutional networks preserve spatial locality, using the hierarchical feature extraction principles established earlier. These properties drive architectural optimizations in AI accelerators, where operations such as data reuse, tiling, and parallel filter computation are important for performance.

The property of translation equivariance is central to understanding why CNNs work effectively for spatial data: shifting the input shifts the output feature map correspondingly. Four aspects connect this property to systems design: the equivariance-invariance distinction, the mathematical formulation, the group theory generalization, and the deployment implications.

Equivariance and invariance are related but distinct concepts that determine how architectures handle transformations. Equivariance means that transforming the input produces the same transformation in the output, as defined in equation 6.4:

$$f(\mathcal{T}(\mathbf{x})) = \mathcal{T}(f(\mathbf{x})) \quad (6.4)$$

For CNNs with translation $\mathcal{T}_v$ (shift by vector v), under stride-1 convolution away from boundary effects (and before any pooling or strided downsampling), if the input shifts by five pixels right, the feature maps also shift by five pixels right. Position information is preserved through the transformation. Invariance, by contrast, means transforming the input does not change the output, as defined in equation 6.5:

$$f(\mathcal{T}(\mathbf{x})) = f(\mathbf{x}) \quad (6.5)$$

Global average pooling over an entire feature map exhibits translation invariance: shifting the input does not change the averaged output. Position information is discarded.

Equivariance matters for learning because it preserves information needed for structured representations. Consider spatial relationships: a feature detector responding to an eye at position (x, y) will respond to the same eye at position $(x + 5, y)$, but the response moves to reflect the new position. The network can learn spatial relationships like “eye above nose” that matter for face detection. Full invariance would lose this relational information, leaving only “eye and nose both present somewhere,” which proves insufficient for many tasks.

Object detection illustrates why equivariance is essential for localization. Detection outputs bounding boxes like “car at (100, 200) with size 50×80 ”, requiring equivariant layers to track position through the network while invariant final layers determine class. This architectural choice matches task structure: equivariance for localization, invariance for classification.

Equivariance also supports hierarchical composition. Early layers detect edges equivariantly at all positions, middle layers combine edges into shapes while maintaining equivariance, and final layers may use partial invariance through pooling for classification. This hierarchy works precisely because intermediate features maintain spatial structure for composition.

These intuitions can be made formal. For a convolutional layer with filter $\mathbf{w}$ and input $\mathbf{x}$, the convolution is

$$(\mathbf{x} * \mathbf{w})[i, j] = \sum_{m, n} \mathbf{w}[m, n] \cdot \mathbf{x}[i + m, j + n].$$

Applying translation $\mathcal{T}_v$ (shift by $v = (v_1, v_2)$) to the input gives $(\mathcal{T}_v \mathbf{x})[i, j] = \mathbf{x}[i - v_1, j - v_2]$. Substituting into the convolution and re-indexing yields

$$((\mathcal{T}_v \mathbf{x}) * \mathbf{w})[i, j] = (\mathbf{x} * \mathbf{w})[i - v_1, j - v_2] = \mathcal{T}_v(\mathbf{x} * \mathbf{w})[i, j],$$

which proves translation equivariance, up to the chosen boundary convention: $f(\mathcal{T}_v \mathbf{x}) = \mathcal{T}_v(f(\mathbf{x}))$.

In practice, the contrast is stark: an equivariant convolutional layer tracks a shifted feature to its new position, preserving the spatial relationships (“whiskers near mouth,” “ears above eyes”) that recognition depends on, while an invariant global pooling layer returns the same scalar wherever the feature appears, discarding position entirely. The worked example that follows traces this tracking through actual matrices.

Example 6.4: Equivariance: Feature detection

Setup: Consider a 7×7 image with a vertical edge at column 3:

$$\mathbf{x} = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Vertical edge detector filter:

$$\mathbf{w} = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

Convolving original image:

Output feature map shows positive activation where the filter transitions from dark to bright (left side of edge) and negative activation where it transitions from bright to dark (right side):

$$f(\mathbf{x}) = \begin{bmatrix} 3 & 0 & -3 & 0 & 0 \\ 3 & 0 & -3 & 0 & 0 \\ 3 & 0 & -3 & 0 & 0 \\ 3 & 0 & -3 & 0 & 0 \\ 3 & 0 & -3 & 0 & 0 \end{bmatrix}$$

Shifted input (edge moved to column 5):

$$\mathcal{T}_2\mathbf{x} = \begin{bmatrix} 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \end{bmatrix}$$

Convolving shifted image:

$$f(\mathcal{T}_2\mathbf{x}) = \begin{bmatrix} 0 & 0 & 3 & 0 & -3 \\ 0 & 0 & 3 & 0 & -3 \\ 0 & 0 & 3 & 0 & -3 \\ 0 & 0 & 3 & 0 & -3 \\ 0 & 0 & 3 & 0 & -3 \end{bmatrix} = \mathcal{T}_2(f(\mathbf{x}))$$

Systems insight: The feature activation shifts by the same amount as the input, demonstrating equivariance. The network knows the edge is at column 5 in the shifted image, not just that an edge exists somewhere.

Equivariance carries systems implications that extend beyond mathematical elegance. Parameter efficiency is the most immediate benefit: equivariance through parameter sharing produces dramatic reductions in model size. Consider processing a 224×224 RGB image. An MLP would require each hidden neuron to connect to all 150,528 input pixels. A CNN with a 3×3 filter needs only 27 parameters per filter, reused across all 224×224 positions. This represents approximately $5,575.1 \times$ fewer parameters per feature detector, and the memory savings enable larger models and bigger batches on fixed hardware.

The computational structure created by equivariance proves equally valuable for systems optimization. The sliding window pattern applies the same operation at every spatial position, creating regular computation that hardware can exploit. Input pixels are used by multiple filter positions, enabling im2col optimizations that restructure data for efficient matrix operations. The resulting computation is inherently SIMD-friendly, as modern GPUs can execute identical instructions across spatial positions simultaneously. This structural regularity explains why TPUs and AI accelerators include specialized units for convolution: the operation maps efficiently to silicon precisely because equivariance creates predictable, parallelizable patterns.

Equivariance also improves sample efficiency in ways that benefit the entire training pipeline. When a network learns an edge detector at one position, equivariance ensures that same detector works at all positions automatically. Training no longer requires examples with edges at every possible location, providing a form of built-in data augmentation. The systems benefits cascade: less training data means reduced storage requirements, faster training, and lower bandwidth consumption during data loading.

Theorem 6.1: Group equivariance formulation

From a group theory perspective, convolution's equivariance to translations represents one instance of a general principle. The translation group $(\mathbb{R}^2, +)$ consists of all 2D translations, closed under composition (translating by v then u equals translating by $v + u$). Convolution is equivariant to this group. Recent research extends this framework to other symmetry groups. Cohen and Welling (2016) developed Group-Equivariant CNNs that handle rotations and reflections by constructing filters equivariant to rotation groups. This allows learning rotation-invariant features for tasks like satellite imagery or medical imaging where orientation does not determine meaning.

The mathematical framework generalizes cleanly: for group G acting on input space X and output space Y , a function $f : X \rightarrow Y$ is G -equivariant if:

$$f(g \cdot \mathbf{x}) = g \cdot f(\mathbf{x}) \quad \forall g \in G, \mathbf{x} \in X$$

Standard CNNs are translation-equivariant, while rotation-equivariant networks extend this to rotation groups. The architectural principle generalizes: data symmetries should be embedded as equivariances in the architecture. For systems engineering, identifying data symmetries directly informs architecture choice: more constrained architectures with stronger symmetries often produce smaller models, and specialized equivariances may require custom operations like rotation convolutions that need either hardware support or efficient software implementations.

In practice, perfect equivariance is often sacrificed for computational efficiency or training stability. Asymmetric padding at image boundaries breaks perfect translation equivariance, as does strided downsampling, which introduces quantization where a one-pixel shift in input produces a noninteger shift in output. Batch normalization, a later normalization layer that stabilizes activations using batch statistics, also breaks equivariance when those statistics are computed per position in some implementations. Modern networks accept these deviations as necessary trade-offs, and the slight loss of theoretical purity rarely impacts practical performance.

Checkpoint 6.2: Spatial inductive bias

CNNs succeed because they match the structure of image data. Verify you understand how:

- ☐ **Parameter sharing:** how does using the same filter across the image reduce parameter count by orders of magnitude compared to MLPs?
- ☐ **Translation equivariance:** why does shifting the input image result in a corresponding shift in the feature map?
- ☐ **Compute-bound convolution:** why does a conv layer reuse each loaded filter weight across many output positions, making it compute-bound compared to other layers?

Different tasks impose different requirements on where equivariance should be maintained vs. where invariance should be introduced. Image classification needs only the final class label to be invariant; intermediate layers benefit from staying equivariant to preserve spatial information for hierarchical feature learning. Object detection requires equivariance throughout the network because bounding box coordinates must track object positions. Semantic segmentation demands full equivariance to the output layer since per-pixel labels must align with input positions. Image generation similarly requires equivariance to maintain spatial structure in the output. The architectural decision of where to introduce invariance through pooling or global averaging vs. maintaining equivariance reflects these task requirements and directly shapes network design.

The preceding task-specific requirements illustrate the inductive bias principle defined in section 6.1: by restricting connectivity to local neighborhoods and sharing parameters across spatial positions, CNNs encode prior knowledge about the structure of visual data—that important features are local and translation-invariant. This architectural constraint reduces the hypothesis space that

the network must search, enabling more efficient learning from limited data compared to fully connected networks.

CNNs naturally implement hierarchical representation learning (Bengio, Courville, et al. 2013) through their layered structure. Early layers detect low-level features like edges and textures with small receptive fields, while deeper layers combine these into increasingly complex patterns with larger receptive fields. This hierarchical organization enables CNNs to build compositional representations: complex objects are represented as compositions of simpler parts. The mathematical foundation for this emerges from stacking convolutional layers, which creates a tree-like dependency structure where each deeper neuron depends on a progressively larger input region; with fixed small kernels, receptive-field side length grows roughly linearly with depth and receptive-field area grows roughly quadratically until it covers the image.

The parameter sharing introduced in section 6.3.1 dramatically reduces complexity compared to MLPs. This sharing embodies the assumption that useful features can appear anywhere in an image, making the same feature detector valuable across all spatial positions.

6.3.3 Computational mapping

How much of this architectural efficiency survives in practice depends on how convolution's sliding-window computation maps onto hardware. Convolution operations create computational patterns distinct from MLP dense matrix multiplication. While high-level frameworks abstract this as a sliding window, the underlying hardware implementation typically transforms the problem to exploit highly optimized matrix multiplication units.

The most common transformation is **im2col** (image-to-column), which rearranges the input image patches into columns of a large matrix, allowing the convolution to be executed as a single GEMM. The computational-primitives discussion uses this transformation to connect CNN structure to matrix hardware.

The bridge between the logical model and physical execution becomes critical for understanding CNN system requirements. While listing 6.3 shows the framework-level abstraction as a simple function call, the hardware must orchestrate complex data movement patterns and exploit spatial locality for efficiency.

Listing 6.3: Convolutional Layer Abstraction: Framework-level convolution operations hide the complexity of sliding window computations, typically dispatching to cuDNN or MKL which internally handle im2col transformations or direct convolution algorithms.

```
def conv_layer_spatial(input, kernel, bias):
    """Framework-level convolution.

    Single call dispatches to optimized kernel (often via im2col + GEMM).
    """
    # Convolution applies shared weights across all positions
    # For a 3x3 kernel on 28x28 input (padded): 9 MACs per position x 784 positions
    output = convolution(input, kernel) + bias
    return activation(output)
```

Listing 6.4 reveals seven nested loops that process each spatial position. While functionally correct, this naive implementation is rarely used in practice due to poor memory locality. Instead, the im2col approach trades memory (duplicating overlapping input pixels) for computational regularity, converting the messy nested loops into a streamlined matrix multiplication that saturates hardware FP units.

The seven nested loops reveal different aspects of the computation. The loop structure divides into three groups: the outer loops manage position, determining which image and where in the image; the middle loop handles output features, computing different learned patterns; and the inner loops perform the actual convolution, sliding the kernel window across the input.

Examining this process in detail, the outer two loops (for y and for x) traverse each spatial position in the output feature map. At each position, values are computed for each output channel (for out_channel loop), representing different learned features or patterns: the 32 different feature detectors.

Listing 6.4: Logical Convolution Computation: Seven nested loops expose $\mathcal{O}(B \times H_{\text{img}} \times W_{\text{img}} \times C_{\text{out}} \times K_h \times K_w \times C_{\text{in}})$ complexity. While logically sound, this pattern is inefficient on hardware; production systems transform this into tiled matrix multiplications.

```
def conv_layer_compute(input, kernel, bias):
    # Logical view of convolution (usually implemented via im2col +
    # GEMM)
    # Loop 1: Process each image in batch
    for image in range(batch_size):
        # Loop 2&3: Move across image spatially
        for y in range(height):
            for x in range(width):
                # Loop 4: Compute each output feature
                for out_channel in range(num_output_channels):
                    result = bias[out_channel]
                    # Loop 5&6: Move across kernel window
                    for ky in range(kernel_height):
                        for kx in range(kernel_width):
                            # Loop 7: Process each input feature
                            for in_channel in range(num_input_channels):
                                # ... MAC operations ...
```

The inner 3 loops implement the actual convolution operation at each position. For each output value, we process a local 3×3 region of the input (the ky and kx loops) across all input channels (for in_channel loop). This creates a sliding window effect, where the same 3×3 filter moves across the image, performing multiply-accumulates between the filter weights and the local input values. Unlike the MLP's global connectivity, this local processing pattern means each output value depends only on a small neighborhood of the input.

With 3×3 filters and 32 output channels, each output position requires only 9 multiply-accumulate operations per input channel, compared to 784 in the reference MLP layer. This operation repeats for every spatial position and every output channel.

While using fewer operations per output, the spatial structure creates different patterns of memory access and computation that systems must handle. These patterns influence system design, creating both challenges and opportunities for optimization. Understanding these system-level implications reveals why CNNs dominate computer vision despite their apparent simplicity.

6.3.4 System implications

The sliding window and im2col transformations described in section 6.3.3 reveal *how* CNNs compute; this section reveals *what that computation costs* in memory, compute, and data movement. These costs depend on layer depth, activation liveness, and whether feature maps remain cached in SRAM or spill to high-bandwidth memory (HBM).

6.3.4.1 Memory requirements

For convolutional layers, memory requirements center around two key components: filter weights and feature maps. Unlike MLPs that require storing full connection matrices, CNNs use small, reusable filters. For a typical CNN processing 224×224 ImageNet images, a convolutional layer with 64 filters of size 3×3 applied to a single input channel requires storing only 576 weight parameters; for multiple input channels, the same kernel and channel product remains dramatically smaller than the millions of weights needed for equivalent fully connected processing. The system must store feature maps for all spatial positions, creating a different memory demand. A 224×224 input with 64 output channels requires storing 3.2M activation values.

These memory access patterns suggest opportunities for optimization through weight reuse and careful feature map management. Processors optimize these spatial patterns by caching filter weights for reuse across positions while streaming feature map data. CPUs use their cache hierarchy to keep frequently used filters resident, while GPUs employ specialized memory architectures designed for the spatial access patterns of image processing. The detailed architecture design principles for these specialized processors are covered in chapter 11.

6.3.4.2 Computation needs

The core computation in CNNs involves repeatedly applying small filters across spatial positions. Each output value requires a local multiply-accumulate operation over the filter region. For ImageNet processing with 3×3 filters and 64 output channels, computing one spatial position involves 576 multiply-accumulates per input channel, and this must be repeated for all 50,176 spatial positions. While each individual computation involves fewer operations than an MLP layer, the total computational load remains large due to spatial repetition.

This computational pattern presents different optimization opportunities than MLPs. The regular, repeated nature of convolution operations enables efficient hardware utilization through structured parallelism. Modern processors exploit this pattern in various ways. CPUs use SIMD instructions¹⁵ to process multiple filter positions simultaneously, while GPUs parallelize computation across spatial positions and channels. The model optimization techniques that further reduce these computational demands, including specialized convolution optimizations and sparsity patterns, are detailed in chapter 10.

6.3.4.3 Data movement

The sliding window pattern of convolutions creates a distinctive data movement profile. Unlike MLPs where each weight is used once per forward pass, CNN filter weights are reused many times as the filter slides across spatial positions. For ImageNet processing, each 3×3 filter weight is reused 50,176 times, once for each position in the 224×224 feature map. The resulting challenge is to stream input features through the computation unit while keeping filter weights stable.

The predictable spatial access pattern enables strategic data movement optimizations. The CPU/GPU caching strategies described earlier apply directly to data movement: frameworks orchestrate computation to maximize 50,176 uses of each filter weight and minimize redundant feature map accesses, exploiting the same spatial locality that makes CNNs memory-efficient.

These memory, compute, and data-movement patterns converge in one model that serves as the chapter’s reference point for compute-bound vision workloads: the ResNet-50 architecture.

 Lighthouse 6.2: ResNet-50 (vision lighthouse)

Why it matters: ResNet-50 is a reference point for regular, convolution-heavy vision workloads. Its architecture consists almost entirely of dense convolutional layers, making it highly regular and efficient on GPUs. Under batched execution with good data reuse, ResNet-50 performance is typically limited by floating-point throughput (FLOP/s), making it a useful lighthouse for explaining data parallelism, quantization, and batching strategies. Table 6.4 summarizes the quantitative properties and their system consequences:

Table 6.4: ResNet-50 Systems Profile: Quantitative properties of the vision lighthouse and their system consequences. Under batched convolutional reuse, ResNet-50 has high effective arithmetic intensity and is typically compute-bound, so peak FP throughput on Tensor Cores often sets the ceiling before memory bandwidth does.

Property	Value	System Implication
Parameters	25.6M	102.4 MB model size at FP32; fits comfortably in GPU memory.
FLOPs/Image	8.2 GFLOP (224×224)	3×3 convolutions are the largest single kernel class at roughly 48% of MACs.
Constraint	Compute-heavy when batched	Limited by peak FLOP/s when weight and activation reuse are high; small-batch inference can move toward the memory-bound regime.
Bottleneck	FP Throughput	Benefits maximally from specialized Matrix Units (Tensor Cores).
Profile	High effective arithmetic intensity under reuse	Arithmetic intensity depends on batch size, convolution algorithm, and materialized memory traffic.

ResNet-50’s compute-bound profile assumes abundant hardware resources, yet most inference runs on devices with power budgets three orders of magnitude smaller than a data center GPU. MobileNetV2 demonstrates that architectural innovation can target this regime, achieving competitive accuracy with a fraction of the computational cost.

¹⁵ | **SIMD (single instruction, multiple data):** CPU instructions that apply the same operation to multiple data elements simultaneously; AVX-512 can process 16 single-precision values per vector instruction; realized speedup over scalar code also depends on instruction throughput, vectorization, and memory access. For CNN inference on edge CPUs without GPU access, SIMD utilization can affect whether a model meets real-time latency targets. Frameworks like TFLite and Open Neural Network Exchange (ONNX) Runtime use vectorized convolution kernels to exploit this parallelism.

Lighthouse 6.3: MobileNetV2 (efficiency lighthouse)

Why it matters: MobileNetV2 represents *latency-constrained* edge workloads. Its depthwise separable convolutions trade channel mixing capacity for speed, making it a useful baseline for mobile apps, embedded vision, and neural architecture search (NAS), automated search over model designs. Table 6.5 summarizes the efficiency lighthouse’s quantitative properties:

Table 6.5: MobileNetV2 Systems Profile: Quantitative properties of the efficiency lighthouse and their system consequences. Roughly an order of magnitude smaller than ResNet-50 in both parameters and FLOPs, MobileNetV2’s operation shapes and arithmetic intensity make single-image latency sensitive to operator-dispatch overhead and memory access on mobile CPUs.

Property	Value	System Implication
Parameters	3.5M	14 MB at FP32; 7.3× smaller than ResNet-50.
FLOPs/Image	600 MFLOP	13.7× fewer than ResNet-50 for similar accuracy.
Constraint	Latency Bound	Single-image inference speed is the priority.
Bottleneck	Overhead/Memory Access	Operator dispatch and memory access can dominate actual compute.
Profile	Low Arithmetic Intensity	Memory access and control logic matter more than peak FLOP/s.

The ResNet-50 and MobileNetV2 profiles create a natural expectation: a model with 13.7× fewer FLOPs should execute proportionally faster. Whether it does depends on operation shapes, implementation, arithmetic intensity, and the target hardware’s roofline balance. MobileNetV2 may achieve less than a proportional speedup on some data center GPUs when its kernels use the available compute units poorly.

With the hardware-mapping caveat in mind, the architectural efficiency of CNNs allows further optimization through specialized techniques like depthwise separable convolutions and pruning [removing low-value weights or channels], detailed in chapter 10. These optimization strategies build on spatial locality principles, with chapter 11 detailing how modern processors exploit convolution’s inherent data reuse patterns.

Systems Perspective 6.1: Misconception: FLOPs equal speed

Fallacy: “MobileNetV2 has 13.7× fewer FLOPs than ResNet-50, so it must run 13.7× faster.”

Resolution: On some high-end GPUs, MobileNetV2 can run slower than ResNet-50 despite using far fewer operations. MobileNetV2’s depthwise separable convolutions have low arithmetic intensity: they move more data relative to computation. GPUs optimized for dense matrix operations may use their compute units poorly on these kernels. FLOPs measure *work*; throughput depends on how well that work *maps to hardware*. This hardware-architecture mismatch recurs as a general fallacy in section 6.11.

6.3.5 Efficient architectures: Keyword spotting

The system implications in section 6.3.4 assume standard CNN architectures with full convolutions. However, standard convolutions scale as $\mathcal{O}(N \times K^2 \times C_{\text{in}} \times C_{\text{out}})$, a cost often prohibitive for the always-on edge devices introduced with our KWS lighthouse. To bridge this gap, efficient architectures like **depthwise separable CNNs (DS-CNN)** decompose the standard convolution into two cheaper operations. This factorization, introduced by Sifre in the context of feature extraction (Sifre and Mallat 2014) and popularized by MobileNet (Howard et al. 2017), reduces cost by separating spatial and channel-wise computation. The depthwise convolution applies filters to each input channel independently ($K \times K \times C_{\text{in}}$ parameters), and the pointwise convolution uses a 1×1 convolution to project channels to the output dimension ($1 \times 1 \times C_{\text{in}} \times C_{\text{out}}$ parameters). This decomposition reduces parameter count and FLOPs to a fraction of roughly $\frac{1}{C_{\text{out}}} + \frac{1}{K^2}$ of standard convolution (yielding a reduction factor of roughly $K^2 \times$, or $\sim 8\text{--}9\times$ for 3×3 filters with large C_{out}), making real-time audio processing feasible on tiny hardware.



Lighthouse 6.4: KWS (TinyML lighthouse)

Why it matters: Keyword spotting models (like DS-CNN) represent the *power-constrained* end of the spectrum. Used in always-on applications like smart doorbells that wake a larger vision pipeline only after an audio trigger, these models must run on microcontrollers with milliwatt power budgets.

KWS forces engineers to count every byte and cycle. It is the lighthouse for extreme quantization (INT8/INT4, detailed in chapter 10) and specialized architectural primitives (Depthwise Separable Convolutions) that trade theoretical representational power for maximum efficiency per watt.

From ResNet-50's compute-heavy standard convolutions through MobileNet's efficient depthwise separable variants to KWS's extreme power-constrained design, CNNs demonstrate how architectural constraints can transform computational challenges into efficiency gains for spatially structured data. Yet their core assumption, that nearby elements are most relevant, fails when patterns depend on temporal order rather than spatial proximity. The next architecture family addresses precisely this limitation.

6.4 RNNs: Sequential Pattern Processing


Convolutional networks exploit *spatial* structure: nearby pixels are more related than distant ones. Many real-world signals, however, have *temporal* structure instead: words in a sentence, samples in an audio stream, sensor readings over time. Processing sequences requires architectures that maintain state across time steps.



Definition 6.4: Recurrent neural networks

Recurrent neural networks (RNNs) are sequence-processing architectures that maintain a hidden state $\mathbf{h}_t = f(\mathbf{h}_{t-1}, \mathbf{x}_t)$ updated at each time step, encoding the assumption that the current output depends on all prior inputs through this fixed-size state vector.

1. **Significance:** The fixed-size state provides $\mathcal{O}(1)$ inference memory regardless of sequence length (processing a 10,000-token sequence requires the same memory as a 10-token sequence), but the sequential update rule creates a sequential bottleneck where all S steps must execute in order, directly contributing to the L_{lat} term of the iron law and making RNNs unable to exploit GPU parallelism across the time dimension during training.
2. **Distinction:** Unlike attention mechanisms, which access the entire token history simultaneously and materialize an $\mathcal{O}(S^2)$ score matrix when processing a full sequence, RNNs compress history into a bottleneck state, meaning gradient signal must propagate back through all S steps—causing $\partial \mathcal{L} / \partial \mathbf{h}_0 \propto \prod_{t=1}^S \partial \mathbf{h}_t / \partial \mathbf{h}_{t-1}$, a product of S Jacobians that vanishes or explodes exponentially with sequence length.
3. **Common pitfall:** A frequent misconception is that RNNs are obsolete. For streaming inference on resource-constrained hardware where $\mathcal{O}(S^2)$ attention memory is prohibitive, such as keyword spotting on a microcontroller, an RNN's $\mathcal{O}(1)$ state size remains the systems-justified choice.



RNNs hold state memory fixed, but latency grows with sequence length.

The limitation manifests concretely in domains such as natural language processing, where word meaning depends on sentential context, and time-series analysis, where future values depend on historical patterns. Sequential data presents a challenge distinct from spatial processing: patterns can span arbitrary temporal distances, rendering fixed-size kernels ineffective. Spatial convolution exploits the principle that nearby pixels are typically related, but temporal relationships operate differently because important connections may span hundreds or thousands of time steps with no correlation to proximity. Traditional feedforward architectures, including CNNs, process each input independently and cannot maintain the temporal context necessary for these long-range dependencies.

Classic recurrent neural networks, from Elman’s simple recurrent network to gated long short-term memory (LSTM) variants, address this architectural limitation (Elman 1990; Hochreiter and Schmidhuber 1997) with a temporal inductive bias: order matters, and the past influences the present. The assumption of sequential dependence guides the introduction of memory as a core component of the computational model. Rather than processing inputs in isolation, RNNs maintain an internal state that propagates information from previous time steps, allowing the network to condition its current output on historical context. This architecture embodies a distinctive trade-off: while CNNs sacrifice theoretical generality for spatial efficiency, recurrent neural networks introduce computational dependencies that challenge parallel execution in exchange for temporal processing capabilities.

6.4.1 Pattern processing needs

Sequential pattern processing addresses scenarios where current input interpretation depends on preceding information. Consider the word “bank”: in “river bank” it denotes a shoreline, but in “bank account” it denotes a financial institution. The correct interpretation depends not just on the word itself but on the words that came before it. This contextual dependency pervades natural language, speech recognition (where phoneme interpretation depends on surrounding sounds), and financial forecasting (where future values depend on historical patterns).

The challenge lies in maintaining and updating relevant context over time. Human text comprehension does not restart with each word; rather, a running understanding evolves as new information arrives. Time-series data compounds this challenge with patterns spanning different timescales, from immediate dependencies to long-term trends. An effective sequential architecture must therefore maintain state over time while updating it in response to new inputs: capturing temporal context in internal state, updating that state as new inputs arrive, and learning which historical information remains relevant for current predictions—all while accommodating variable-length sequences that MLPs and CNNs cannot naturally handle.

6.4.2 Algorithmic structure

Section 6.4.1 demands an architecture that maintains and updates state over time. RNNs address this through recurrent connections, distinguishing them from MLPs and CNNs. Rather than merely mapping inputs to outputs, RNNs maintain an internal state updated at each time step, creating a memory mechanism that propagates information forward in time. This temporal dependency modeling capability was demonstrated influentially by Elman (1990), whose experiments identified structure in time-dependent data. Basic RNNs suffer from the vanishing gradient problem, constraining their ability to learn long-term dependencies.

Equation 6.6 updates the hidden state from the current input and preceding state.

$$\mathbf{h}_t = f(\mathbf{W}_{\text{hh}}\mathbf{h}_{t-1} + \mathbf{W}_{\text{hx}}\mathbf{x}_t + \mathbf{b}_h) \quad (6.6)$$

where $\mathbf{h}_t$ denotes the hidden state at time t , $\mathbf{x}_t$ denotes the input at time t , $\mathbf{W}_{\text{hh}}$ contains the recurrent weights, $\mathbf{W}_{\text{hx}}$ contains the input weights, $\mathbf{b}_h$ is the hidden-state bias vector, and f is the activation function. The equation uses column vectors; the later listings use row-major batched tensors, so their weight multiplications appear on the right. Compare the left and right panels of figure 6.5: the left panel shows the compact recurrent loop, while the right panel unfolds it across time steps, making explicit the temporal dependencies that this recurrence creates.

In word sequence processing, each word may be represented as a 100-dimensional vector ($\mathbf{x}_t$), with a hidden state of 128 dimensions ($\mathbf{h}_t$). At each time step, the network combines the current input with its previous state to update its sequential understanding, establishing a memory mechanism capable of capturing patterns across time steps.

This recurrent structure fulfills sequential processing requirements through connections that maintain internal state and propagate information forward in time. Rather than processing all inputs independently, RNNs process sequential data by iteratively updating a **hidden state** based on the current input and the previous hidden state. This architecture suits tasks including language modeling, speech recognition, and time-series forecasting.

RNNs implement a recursive algorithm where each time step’s function call depends on the result of the previous call. Analogous to recursive functions that maintain state through the call stack, RNNs




Figure 6.5: Recurrent Neural Network Unfolding: Left panel shows the compact recurrent loop; right panel unfolds this across time steps. Three weight matrices are shared across all steps: W_{hx} (input-to-hidden), W_{hh} (hidden-to-hidden), and W_{yh} (hidden-to-output). These matrices are reused at every step, so the parameter count is independent of sequence length. For a 128-dimensional hidden state and a 100-dimensional input, each time step requires 16,384 MACs for recurrent connections plus 12,800 for input projection.

maintain state through their hidden vectors. The mathematical formula $\mathbf{h}_t = f(\mathbf{h}_{t-1}, \mathbf{x}_t)$ directly parallels recursive function definitions where $f(n) = g(f(n-1), \text{input}(n))$. This correspondence explains RNN capacity to handle variable-length sequences: just as recursive algorithms process lists of arbitrary length by applying the same function recursively, RNNs process sequences by applying the same recurrent computation. This sequential dependency has a direct hardware consequence. The recurrence imposes a dependency at every time step because $\mathbf{h}_t$ cannot begin until $\mathbf{h}_{t-1}$ is available. The resulting $\mathcal{O}(S)$ critical path limits parallelism across time steps. Actual accelerator utilization depends on hidden dimensions, batch size, kernel implementation, sequence length, and hardware.

Sequential processing creates computational bottlenecks but produces unique efficiency characteristics for memory usage. RNNs' $\mathcal{O}(d_{\text{hidden}})$ inference memory overhead (analyzed in detail in section 6.4.4) stays flat in sequence length, where an autoregressive transformer's key-value cache grows as $\mathcal{O}(S \cdot d_{\text{model}})$, allowing processing of sequences thousands of steps long on modest hardware. During training with backpropagation through time (BPTT), however, RNNs must store activations for all time steps, requiring $\mathcal{O}(S \cdot d_{\text{hidden}})$ memory. The recurrent weight matrix often contains connections with minimal contribution to temporal dependencies, allowing significant compression through methods covered in chapter 10.

6.4.3 Computational mapping

RNN sequential processing creates computational patterns different from both MLPs and CNNs, extending the architectural diversity discussed in section 6.1. This implementation approach shows temporal dependencies translating into specific computational requirements.

Listing 6.5 demonstrates the single-time-step mechanism using framework-level matrix operations: combine the previous hidden state with the current input, add the bias, and apply the activation to produce the next hidden state. The code is intentionally local to one step because the systems cost is not the step itself, but the dependency chain that prevents parallel execution across time.

Listing 6.5: RNN Layer Abstraction: Framework-level implementation combining two matrix multiplications, $\mathbf{h}_{t-1} W_{hh}$ and $\mathbf{x}_t W_{hx}$, per time step. For 128-dimensional hidden state and 100-dimensional input, each step requires $16,384 + 12,800 = 29,184$ MACs, but sequential dependencies prevent parallelization across time.

```
def rnn_layer_step(x_t, h_prev, W_hh, W_hx, b):
    # x_t: input at time t (batch_size × input_dim)
    # h_prev: previous hidden state (batch_size × hidden_dim)
    # W_hh: recurrent weights (hidden_dim × hidden_dim)
    # W_hx: input weights (input_dim × hidden_dim)
    h_t = activation(matmul(h_prev, W_hh) + matmul(x_t, W_hx) + b)
    return h_t
```


The function handles a single time step, taking the current input $\mathbf{x}_t$ and previous hidden state $\mathbf{h}_{\text{prev}}$, along with two weight matrices: $\mathbf{W}_{\text{hh}}$ for hidden-to-hidden connections and $\mathbf{W}_{\text{hx}}$ for input-to-hidden connections. Through matrix multiplication operations (`matmul`), it merges the previous state and current input to generate the next hidden state.

The simple recurrence $\mathbf{h}_t = \tanh(\mathbf{W}_{\text{hh}}\mathbf{h}_{t-1} + \mathbf{W}_{\text{hx}}\mathbf{x}_t + \mathbf{b})$ conceals a computational structure with unique challenges: sequential dependencies that prevent parallelization, memory access patterns that differ from feedforward networks, and state management requirements that affect system design. The detailed implementation in listing 6.6 reveals the computational reality beneath the mathematical abstraction. Its nested loop structure exposes how sequential processing creates both limitations and opportunities in system optimization.

Listing 6.6: RNN Layer Computation: Nested loops expose the sequential dependency structure. Loop 1 enables batch parallelism, but Loops 2-3 must complete before Loop 4's activation, and crucially, each time step depends on the previous step's output, creating $\mathcal{O}(S)$ sequential depth that prevents GPU parallelization across the time dimension.

```
def rnn_layer_compute(x_t, h_prev, W_hh, W_hx, b):
    # Initialize next hidden state
    h_t = np.zeros_like(h_prev)

    # Loop 1: Process each sequence in the batch
    for batch in range(batch_size):
        # Loop 2: Compute recurrent contribution (h_prev * W_hh)
        for i in range(hidden_dim):
            for j in range(hidden_dim):
                h_t[batch, i] += h_prev[batch, j] * W_hh[j, i]

        # Loop 3: Compute input contribution (x_t * W_hx)
        for i in range(hidden_dim):
            for j in range(input_dim):
                h_t[batch, i] += x_t[batch, j] * W_hx[j, i]

        # Loop 4: Add bias and apply activation
        for i in range(hidden_dim):
            h_t[batch, i] = activation(h_t[batch, i] + b[i])

    return h_t
```

The nested loops in `rnn_layer_compute` expose the core computational pattern of RNNs. Loop one processes each sequence in the batch independently, allowing for batch-level parallelism. Within each batch item, Loop two computes how the previous hidden state influences the next state through the recurrent weights $\mathbf{W}_{\text{hh}}$. Loop three then incorporates new information from the current input through the input weights $\mathbf{W}_{\text{hx}}$. Finally, Loop four adds biases and applies the activation function to produce the new hidden state.

For a sequence processing task with input dimension 100 and hidden state dimension 128, each time step requires two matrix multiplications: one 128×128 for the recurrent connection and one 100×128 for the input projection. While individual time steps can process in parallel across batch elements, the time steps themselves must execute sequentially, producing a computational pattern with fundamentally different parallelization characteristics than MLPs or CNNs.

6.4.4 System implications

RNNs introduce an inescapable system constraint: *sequential dependency*. Unlike MLPs and CNNs where parallelism scales with the number of neurons or pixels, RNN parallelism is limited across the sequence dimension. Increasing peak compute rate or memory bandwidth can accelerate each step, but it cannot remove the dependency chain along the sequential critical path.

The core computation $\mathbf{h}_t = \tanh(\mathbf{W}_{\text{hh}}\mathbf{h}_{t-1} + \mathbf{W}_{\text{hx}}\mathbf{x}_t)$ creates a strict ordering. Time step t cannot begin until step $t-1$ completes. If processing a document with 1,000 words, the system must execute 1,000 dependent matrix-vector multiplications. Additional hardware can reduce the latency of each multiplication but cannot parallelize the dependent time steps. This limits the parallel width to the batch size, whereas CNNs can exploit parallelism across spatial dimensions, channels, and batches.

RNNs are uniquely memory-efficient for long sequences during inference. They maintain a fixed-size hidden state vector (for example, 2 KB for a 512-dim state) regardless of whether the sequence length is 10 or 10,000. This $\mathcal{O}(d_{\text{hidden}})$ state scaling contrasts with attention, introduced next, which retains sequence state that grows with length: full transformer attention during training or prompt processing stores score interactions that scale as $\mathcal{O}(S^2)$, while autoregressive transformer serving keeps an $\mathcal{O}(Sd_{\text{model}})$ key-value cache. The compression comes at a cost, however: the fixed-size state becomes an information bottleneck, forcing the network to compress arbitrary history into a small vector and leading to the vanishing gradient problems that motivated LSTMs and eventually transformers.

RNNs exhibit high temporal locality for weights (reused every step) but low locality for activations. The weight matrices $\mathbf{W}_{\text{hh}}$ and $\mathbf{W}_{\text{hx}}$ stay in the cache (or on-chip memory) for the entire duration of the sequence processing, achieving high arithmetic intensity if the batch size is large enough. However, the requirement to read and write the hidden state at every step creates a constant stream of low-intensity updates that can strain memory bandwidth if not carefully managed.

This tension between memory efficiency and sequential execution defined the pre-transformer era. RNNs compress arbitrarily long histories into a fixed-size hidden state, which is memory efficient but creates two compounding problems: the sequential dependency prevents hardware from parallelizing across time steps, and the fixed-capacity state becomes an information bottleneck where early inputs fade as sequences grow (the vanishing gradient problem). Together, these limitations motivated architectures that could access any position in a sequence directly, without processing all intervening elements. That direct-access capability, developed in section 6.5, is the attention mechanism. Hardware strategies for managing sequential bottlenecks in RNN workloads that remain in production, including pipeline parallelism and operator fusion, are analyzed in section 11.8.

6.5 Attention: Dynamic Processing

The RNN bottlenecks analyzed in section 6.4.4 become concrete with a simple example. Consider the sentence “The cat, which was sitting by the window overlooking the garden, was sleeping.” Here, “cat” and “sleeping” are separated by multiple intervening words, yet they form the core subject-predicate relationship. An RNN would process all intervening elements sequentially, potentially losing this connection in its fixed-capacity hidden state. This limitation motivates an alternative: an architecture that directly computes the relevance between any two positions regardless of distance.

Attention mechanisms¹⁶ address precisely this challenge (Bahdanau et al. 2015) by introducing dynamic connectivity patterns that adapt based on input content. Rather than relying on a single fixed-length representation, the original encoder-decoder attention mechanism computes relevance between each decoder state and the encoder’s source positions and weights those interactions accordingly.



Definition 6.5: Attention mechanisms

Attention mechanisms are neural network operations that compute a weighted sum of value vectors, where the weights are derived from learned similarity scores between a query vector and a set of key vectors, enabling dynamic, content-dependent information routing between any two positions in a sequence.

1. **Significance:** Attention connects any two tokens in $\mathcal{O}(1)$ depth but computes S^2 score interactions. If those scores are materialized, a 4,096-token sequence with 16-bit scores consumes 33.6 MB per layer per head (about 16.8M scores at 2 bytes each), directly increasing the iron law’s D_{vol} and BW terms. Tiled exact-attention kernels avoid retaining the full matrix, but the dense score computation remains quadratic.
2. **Distinction:** Unlike RNNs, which compress all prior context into a single fixed-size state vector, attention mechanisms retain token representations and compute relevance scores directly. During training or full-sequence prefill, the initial pass that processes the whole prompt, score interactions grow quadratically with sequence length; during

¹⁶ | **Bahdanau attention:** This approach broke the “fixed-length vector” bottleneck of prior sequence-to-sequence models by allowing a decoder to dynamically query all input elements at each output step, creating the adaptive connectivity described. This replaced the structural constraint of a fixed-capacity channel with a learned, content-based weighting system. The core trade-off was accepting a linear, $\mathcal{O}(S)$ memory cost to store all input states in exchange for overcoming the information loss inherent in a single vector.

autoregressive serving, the stored KV cache, the saved key and value vectors from prior tokens, grows as $\mathcal{O}(Sd_{\text{model}})$ while each new token attends over prior keys and values.

3. **Common pitfall:** A frequent misconception is that attention is a general-purpose weighting scheme that can be applied freely. Quadratic score computation is a hard scaling constraint: doubling the context quadruples score interactions. Naive implementations also quadruple score storage; tiled exact algorithms such as FlashAttention avoid the full matrix, while sparse variants reduce the scores computed.

While attention mechanisms were initially used as components within recurrent architectures, their ability to connect any position to any other made the recurrent structure unnecessary for many sequence tasks. The transformer¹⁷ architecture (Vaswani et al. 2017) combined attention with feed-forward layers, residual connections, normalization, and positional information without recurrence. This architectural shift traded the RNN's $\mathcal{O}(S)$ sequential path for constant path length between positions within one attention layer, enabling parallelization across sequence positions on high-throughput accelerators.

The transformer architecture in section 6.6 inherits every important systems property from attention itself: dynamic routing, parallel sequence processing, and quadratic score construction. The next step is therefore to establish what kind of pattern-processing problem attention solves before treating the transformer as a full architecture.

6.5.1 Pattern processing needs

Dynamic pattern processing addresses scenarios where relationships between elements are not fixed by architecture but instead emerge from content. Language translation exemplifies this challenge: when translating “the bank by the river,” understanding “bank” requires attending to “river,” but in “the bank approved the loan,” the important relationship is with “approved” and “loan.” Unlike RNNs that process information sequentially or CNNs that use fixed spatial patterns, an architecture is required that can dynamically determine which relationships matter. The pronoun-resolution schematic in figure 6.6 makes that dynamic routing visible.

The student did not finish all the homework because they were tired.

Layer: 4 Head: 2




Figure 6.6: Illustrative Attention Pattern: A schematic attention pattern for the pronoun “they.” Line thickness represents hypothetical attention-weight magnitude, with stronger links to “student,” “finish,” and “homework.” One attention layer connects all positions directly and can compute their dense interactions in parallel.

Input-dependent processing extends well beyond language. In protein structure prediction, amino-acid interactions depend on chemical properties and spatial arrangement, not only chain position. In

¹⁷ | **Transformer:** The founding paper, “Attention Is All You Need,” made the explicit systems claim that a parallel attention mechanism could fully replace sequential recurrent processing. This architectural trade eliminates an RNN’s $\mathcal{O}(S)$ path length constraint on parallelism but introduces $\mathcal{O}(S^2)$ dense score computation; naïve implementations also materialize $\mathcal{O}(S^2)$ score storage. This quadratic computation continues to shape context-window engineering.

graph analysis, graph convolutional networks (GCNs) aggregate features over each node’s observed neighbors (Kipf and Welling 2017). Unlike CNNs, which access memory in regular spatial strides, or transformers, which work with dense sequence tensors, GCN neighbor aggregation follows the irregular adjacency structure of the input graph. Each gather step touches an input-specific set of node embeddings, defeating cache prefetchers and preventing coalesced memory access. Irregular neighbor-gather patterns can therefore bind execution on memory bandwidth and cache misses rather than compute throughput. In document analysis, connections between sections depend on semantic content rather than proximity.

Input dependence does not imply all-to-all connectivity. Dense attention scores every pair of sequence elements, whereas a GCN restricts aggregation to edges in the observed graph. Input-specific graphs can require different gather schedules and exhibit different locality. The common systems challenge is that the computation or memory-access pattern responds to the input rather than following a fixed stencil. Even when tensor dimensions stay fixed, content-dependent scores determine which information is emphasized. Attention provides one mechanism for learning such relationships, and dense self-attention forms the foundation of the transformer architecture.

When processing the pronoun “they” in the sentence, an attention mechanism can assign larger weights to context such as “student” and “finish,” as the schematic line thickness illustrates. This direct interaction can represent long-range dependencies without recurrently processing every intervening token.

6.5.2 Algorithmic structure

Section 6.5.1 requires computing relationships dynamically based on content. Attention mechanisms achieve this by computing weighted connections between elements based on their content (Bahdanau et al. 2015), processing relationships that emerge from the data itself rather than being fixed by architecture. Transformers use the scaled dot-product form introduced by Vaswani et al. (2017):

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}} \right) \mathbf{V}$$

This equation shows scaled dot-product attention. $\mathbf{Q}$ (queries) and $\mathbf{K}$ (keys) are matrix-multiplied to compute similarity scores using the dot product; section C.1.3 formalizes the dot product as a similarity measure. The scores are divided by $\sqrt{d_k}$ (key dimension) for numerical stability, then normalized with softmax¹⁸ to produce attention weights. These weights are applied to $\mathbf{V}$ (values) to produce the output. The result is a weighted combination where each position receives information from all relevant positions based on content similarity.

In this equation, $\mathbf{Q}$ (queries), $\mathbf{K}$ (keys), and $\mathbf{V}$ (values)¹⁹ represent learned projections of the input. For a sequence of length S with dimension d , this operation creates an $S \times S$ attention matrix, determining how each position should attend to all others.

The attention operation involves several key steps. First, it computes query, key, and value projections for each position in the sequence. In figure 6.7, each cell in the $S \times S$ attention matrix represents a query-key interaction, and in a trained model the largest entries mark which positions attend most strongly to which others. Finally, these attention weights combine value vectors to produce the output.

Unlike the fixed weight matrices found in previous architectures, attention weights are computed dynamically for each input. Follow the matrix dimensions in figure 6.8 to see this dynamic computation unfold: the embedding matrix multiplies with QKV weight matrices in a single batched operation, and the resulting projections change for every new input sequence.

6.5.3 Computational mapping

Self-attention scales quadratically with sequence length S , creating severe HBM memory pressure for long inputs. Listing 6.7 takes the projected queries, keys, and values and counts the multiply-accumulate operations that the all-pairs score matrix and the value aggregation each require.

The translation from attention’s mathematical elegance to hardware execution reveals the computational price of dynamic connectivity. While the attention equation $\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}(\mathbf{Q}\mathbf{K}^T / \sqrt{d_k})\mathbf{V}$ appears as a straightforward matrix operation, the physical implementation

¹⁸ | **Softmax:** Named as a “soft” (differentiable) version of argmax, with mathematical roots in Boltzmann’s statistical mechanics, softmax normalizes each query row over all S scores; a naive attention implementation materializes the full $S \times S$ matrix before normalization. Online softmax and tiling, as used by FlashAttention, instead maintain running normalization statistics and compute exact attention without retaining that full matrix in HBM. Softmax therefore requires a reduction over each row, but it does not force quadratic auxiliary storage; the dense pairwise score computation remains quadratic.

¹⁹ | **Query-key-value (QKV):** The terminology is borrowed from information retrieval, explaining why the equation uses three distinct learned projections to calculate pairwise scores. Creating these projections requires three independent weight matrices, costing $3 \times d_{\text{model}}^2$ parameters per layer. The direct systems consequence is the “KV cache” for autoregressive inference: all prior key and value vectors must be stored for the next token, causing memory to grow linearly ($\mathcal{O}(S)$) with sequence length and potentially dominate serving memory at long contexts or high concurrency.




Figure 6.7: Query-Key-Value Attention Mechanism: For a 6-token sequence, queries (cyan) match against keys (red) to produce a 6×6 attention matrix with $\mathcal{O}(S^2)$ entries. Cell color runs from blue to red across the grid as an illustrative gradient rather than as computed weights. Each output position aggregates information from all values (green) weighted by its own row of that matrix. The matrix structure reveals both the computational pattern (thirty-six similarity computations) and the naive memory bottleneck (storing S^2 attention weights). Source: Transformer Explainer (Cho et al. 2025).




Figure 6.8: QKV Projection Computation: The embedding matrix (6×768) multiplies with QKV weight matrices (768×2304) plus bias to produce combined projections (6×2304). The 2304 output dimension contains concatenated query, key, and value projections (each 768-dimensional). This single batched matrix multiplication, requiring $6 \times 768 \times 2304 = 10.6$ million MACs, replaces 3 separate projection operations for efficiency. Source: Transformer Explainer (Cho et al. 2025).

requires orchestrating quadratic numbers of pairwise computations that create different system demands than previous architectures. The nested loops in `attention_layer_compute` expose this computational signature. The first loop processes each sequence in the batch independently. The second and third loops compute attention scores between all pairs of positions, creating the quadratic computation pattern that makes attention both powerful and computationally demanding. The fourth loop uses these attention weights to combine values from all positions, completing the dynamic connectivity pattern that defines attention mechanisms.

Listing 6.7: Attention Computation: Two implementations showing the same $\mathcal{O}(S^2 \times d)$ complexity. The matrix form (top) uses optimized GEMM, while the nested loops (bottom) expose the quadratic pairwise comparisons: for sequence length 512 and dimension 64, computing attention scores requires $512 \times 512 \times 64 = 16.8\text{M}$ MACs per attention head, plus another 16.8M MACs for value aggregation.

```
def attention_layer_matrix(Q, K, V):
    # Q, K, V: (batch_size × seq_len × d_model)
    # Compute attention scores
    scores = matmul(Q, K.transpose(-2, -1)) / sqrt(d_k)
    weights = softmax(scores) # Normalize scores
    output = matmul(weights, V) # Combine values
    return output

def attention_layer_compute(Q, K, V):
    # Initialize outputs
    scores = np.zeros((batch_size, seq_len, seq_len))
    outputs = np.zeros_like(V)

    # Loop 1: Process each sequence in batch
    for b in range(batch_size):
        # Loop 2: Compute attention for each query position
        for i in range(seq_len):
            # Loop 3: Compare with each key position
            for j in range(seq_len):
                # Compute attention score
                for d in range(d_model):
                    scores[b, i, j] += Q[b, i, d] * K[b, j, d]
                scores[b, i, j] /= sqrt(d_k)

            # Apply softmax to scores
            for i in range(seq_len):
                scores[b, i] = softmax(scores[b, i])

        # Loop 4: Combine values using attention weights
        for i in range(seq_len):
            for j in range(seq_len):
                for d in range(d_model):
                    outputs[b, i, d] += scores[b, i, j] * V[b, j, d]


    return outputs
```

6.5.4 System implications

Attention mechanisms exhibit distinctive system-level patterns that differ from previous architectures through their dynamic connectivity requirements. In iron law terms (section 1.7), attention shifts the bottleneck from the latency-bound sequential path of RNNs toward quadratic score interactions and implementation-dependent data movement. Naive implementations materialize an $\mathcal{O}(S^2)$ attention matrix, while tiled algorithms such as FlashAttention avoid storing the full matrix by recomputing and streaming blocks through faster memory (see section C.5 for the step-by-step online softmax recurrence and I/O reduction derivation).

6.5.4.1 Memory requirements

Attention mechanisms require storage for query-key-value projections and intermediate feature representations. A naive implementation also materializes an $S \times S$ attention-weight matrix for every sequence and head, creating a quadratic memory bottleneck alongside the $S \times d$ inputs and outputs. For long sequences, these transient scores can dominate accelerator memory even though the learned projection matrices remain fixed in size. Tiled exact-attention kernels such as FlashAttention stream score blocks through faster memory and avoid retaining the complete matrix in HBM, while still computing S^2 dense interactions. This distinction separates an unavoidable quadratic compute cost for dense attention from optional quadratic score storage. A quick calculation shows how fast the naive storage wall appears at scale.



Doubling context quadruples score interactions and naive score storage.

Napkin Math 6.1: The quadratic bottleneck

Problem: How much memory does a materialized attention matrix require at sequence length $S = 100,000$ (context window)?

Math:

1. **Matrix size:** The attention score matrix ($\mathbf{QK}^T$) has dimensions $S \times S$ per head, across N_{heads} heads (12 heads in this example).
2. **Elements:** $100,000 \times 100,000 \times 12 \text{ heads} = 1.2 \times 10^{11}$ elements.
3. **Memory:** At FP16 (2 bytes/element): $1.2 \times 10^{11} \times 2 \text{ bytes} = 240 \text{ GB}$.
4. **Retained layers:** $240 \text{ GB per layer} \times 32 \text{ retained layers} = 7,680 \text{ GB}$.

Systems insight: A materialized attention matrix consumes 240 GB per layer. Naive training that retains all 32 layers' scores for backward would require 7,680 GB, far exceeding any single GPU's capacity. This memory wall motivates two broad implementation strategies developed later: avoid materializing the full matrix by tiling the computation, or reduce the number of scores computed in the first place.

6.5.4.2 Computation and data movement

Attention computation divides into two main phases: generating attention weights and applying them to values. For each attention layer, the system performs many multiply-accumulate operations across multiple computational stages. The query-key interactions alone require $S \times S \times d$ multiply-accumulates, with an equal number needed for applying attention weights to values. Additional computations are required for the projection matrices and softmax operations. This computational pattern differs from previous architectures due to its quadratic scaling with sequence length and the need to perform fresh computations for each input.

Checkpoint 6.3: Quadratic scaling intuition

Long-context scaling is shaped by the cost of attention. Verify your intuition:

- ☐ **Quadratic growth:** can you explain why doubling sequence length quadruples dense score interactions and a materialized score matrix?
- ☐ **Tiling boundary:** can you explain why long-context dense attention remains more expensive even when tiling removes full-matrix storage?

Data movement in attention mechanisms presents challenges distinct from all previous architectures. Each attention operation requires projecting and moving query, key, and value vectors for every position, then coordinating value movement during the weighted combination. A naive implementation also stores and accesses the full $S \times S$ attention matrix, making intermediate scores a major bandwidth cost; tiled exact attention changes this traffic pattern. Unlike the predictable spatial access patterns of CNNs or the sequential access of RNNs, attention moves dynamically computed scores and vectors across the memory hierarchy, complicating simple caching strategies.

Example 6.5: The quadratic wall

Scenario: Early BERT deployments enforced rigid 512-token context sequence limits during inference and prefill (Devlin et al. 2019).

Diagnosis: Standard self-attention materializes an $S \times S$ score matrix scaling quadratically ($\mathcal{O}(S^2)$) in memory. Increasing sequence length from 512 to 4,096 increases score memory by 64×, triggering GPU out-of-memory crashes.

Systems lesson: Super-linear memory scaling turns algorithmic complexity into hardware bottlenecks. Serving long sequences requires FlashAttention IO-tiling, sparse attention, or strict token truncation to bound SRAM activation footprint.

The memory, computation, and data movement characteristics of attention shape system design and raise the question of whether recurrence is still necessary. Early transformer deployments commonly enforced short context windows because materialized attention scores could exhaust device memory. Despite this cost, attention connects any two positions in constant depth and can replace the sequential processing path of recurrent architectures. This trade-off motivated the transformer architecture.

6.6 Transformers: Parallel Sequence Processing

Attention provides the computational primitive of dynamic, content-dependent routing between positions, yet it was originally layered on top of recurrent architectures, inheriting their sequential bottleneck. The **transformer architecture** removes recurrence by combining attention with feed-forward layers, residual connections, normalization, and positional information. This enables parallel computation across sequence positions during full-sequence processing while retaining dynamic connectivity. The trade also creates important system costs: quadratic dense-score computation, growing key-value state during serving, and high bandwidth pressure during some autoregressive-generation regimes.



Definition 6.6: Transformers

Transformers are neural network architectures that combine self-attention, feed-forward layers, residual connections, normalization, and positional information without recurrence, enabling parallel processing across sequence positions during training and prefill.

1. **Significance:** Parallelism is the systems payoff. An RNN processing an S -token sequence executes S dependent steps. A transformer can process the positions of a full sequence in parallel, and one attention layer provides constant path length between any two positions. Matrix utilization depends on sequence length, batch size, model dimensions, implementation, datatype, and hardware. Dense attention computes $\mathcal{O}(S^2)$ score interactions; naive implementations also materialize $\mathcal{O}(S^2)$ score storage.
2. **Distinction:** Unlike attention used inside a recurrent backbone, which inherits the host architecture's $\mathcal{O}(S)$ sequential path, transformers use attention as the primary mixing operation between positions and combine it with position-wise feed-forward layers and supporting components.
3. **Common pitfall:** A frequent misconception is that transformers have “infinite context.” Context length is bounded by two distinct memory costs: the attention-score matrix during training and prefill, and the KV cache that accumulates during autoregressive serving. At long contexts the cache alone can rival the model weights, which is why KV cache compression and related serving optimizations are active engineering areas rather than optional refinements.

6.6.1 Pattern processing needs

Attention mechanisms first appeared as additions to existing architectures, particularly RNN-based sequence-to-sequence tasks (Sutskever et al. 2014; Bahdanau et al. 2015). Those hybrids improved dynamic connectivity but kept the recurrent bottleneck: limited parallelism and difficulty with very long sequences. The transformer section begins from the architectural decision to remove that bottleneck entirely, then follows the cost of that decision through memory, bandwidth, and serving state.

Transformers, introduced in the “Attention Is All You Need” paper by Vaswani et al. (2017), embody a different inductive bias: self-attention permits all-to-all interaction, while positional mechanisms encode sequence order. Rather than adding attention to RNNs, transformers built the architecture around self-attention as the primary mixing operation. This architectural decision traded the parameter efficiency of CNNs and the recurrence of RNNs for flexible interactions and parallel processing of full sequences.

6.6.2 Algorithmic structure

The key innovation in transformers lies in their use of self-attention layers. In the self-attention mechanism used by transformers, the query, key, and value vectors are all derived from the same input sequence. This is the key distinction from earlier attention mechanisms where the query might come from a decoder while the keys and values came from an encoder. By making all components self-referential, self-attention allows the model to weigh the importance of different positions within the same sequence when encoding each position. For instance, in processing the sentence “The animal did not cross the street because it was too wide,” self-attention allows the model to link “it” with “street,” capturing long-range dependencies that are challenging for traditional sequential models.

The self-attention mechanism differs from earlier attention in one critical respect: every query, key, and value is derived from the same input $\mathbf{X}$, as equation 6.7 makes explicit:

$$\text{SelfAttention}(\mathbf{X}) = \text{softmax} \left(\frac{\mathbf{XW}_Q (\mathbf{XW}_K)^T}{\sqrt{d_k}} \right) \mathbf{XW}_V \quad (6.7)$$

Here, $\mathbf{X}$ is the input sequence, and $\mathbf{W}_Q$, $\mathbf{W}_K$, and $\mathbf{W}_V$ are learned weight matrices for queries, keys, and values respectively. This formulation highlights how self-attention derives all its components from the same input, creating a dynamic, content-dependent processing pattern.

Building on this foundation, transformers employ multi-head attention, which extends the self-attention mechanism by running multiple attention functions in parallel. Each “head” involves a separate set of query/key/value projections that can focus on different aspects of the input, allowing the model to jointly attend to information from different representation subspaces. This multi-head structure provides the model with a richer representational capability, enabling it to capture various types of relationships within the data simultaneously.

Each head learns a separate projection into its own subspace, and their outputs are concatenated and linearly mixed, as equation 6.8 formalizes:

$$\text{MultiHead}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Concat}(\text{head}_1, \dots, \text{head}_{N_{\text{heads}}}) \mathbf{W}^O \quad (6.8)$$

where each attention head is computed as:

$$\text{head}_i = \text{Attention}(\mathbf{QW}_i^Q, \mathbf{KW}_i^K, \mathbf{VW}_i^V)$$

A critical component in both self-attention and multi-head attention is the scaling factor $\sqrt{d_k}$, which serves an important mathematical purpose. This factor prevents the dot products from growing too large, which would push the softmax function into regions with extremely small gradients. If the query and key components are independent, zero mean, and unit variance, their dot product has variance d_k , so dividing by $\sqrt{d_k}$ normalizes the variance to one, maintaining stable gradients and enabling effective learning.²⁰

Beyond the mathematical mechanics, attention mechanisms can be understood conceptually as implementing a form of content-addressable memory system. Like hash tables that retrieve values based on key matching, attention computes similarity between a query and all available keys, then retrieves a weighted combination of corresponding values. The dot product similarity $\mathbf{q}_i \cdot \mathbf{k}_j$ functions like a hash function that measures how well each key matches the query. The softmax normalization ensures the weights sum to one, implementing a probabilistic retrieval mechanism. This connection explains why attention proves effective for tasks requiring flexible information retrieval: it provides a differentiable approximation to database lookup operations.

From an information-theoretic perspective, attention mechanisms implement smooth information aggregation. The softmax weights form a probability distribution over keys, and the distribution’s entropy measures how concentrated or diffuse those weights are (Cover and Thomas 2006). This smooth weighting lets the mechanism combine information from several positions rather than making a hard retrieval decision.

Attention mechanisms exhibit significant redundancy (many heads learning similar patterns), and the softmax operation creates sensitivity to reduced precision. These properties create opportunities for optimization through pruning, factorization, sparse attention patterns, and specialized quantization, all covered in chapter 10.

²⁰ | **Attention scaling** ($\sqrt{d_k}$): This normalization directly counteracts the linear growth in variance (d_k) of the query-key dot product, preventing the softmax function from saturating where gradients would otherwise vanish. The systems consequence is most acute in mixed-precision training: when activations grow large, unscaled dot products can produce logits that overflow the 16-bit float range, destabilizing or halting learning entirely.

This information-theoretic interpretation reveals why attention is effective for selective processing. The mechanism balances two competing objectives: focusing probability mass on high-scoring positions while preserving enough entropy to avoid a brittle hard selection. This smooth weighting helps transformers handle long sequences and complex dependencies.

Self-attention learns dynamic activation patterns across the input sequence. Unlike CNNs which apply fixed filters or RNNs which use fixed recurrence patterns, attention learns which elements should activate together based on their content. This creates a form of adaptive connectivity where the effective network topology changes for each input. Recent research has shown that attention heads in trained models often specialize in detecting specific linguistic or semantic patterns (Clark et al. 2019), suggesting that the mechanism naturally discovers interpretable structural regularities in data.

The transformer architecture applies this self-attention mechanism within a broader structure that typically includes feed-forward layers, layer normalization, and residual connections. Figure 6.9 shows input tokens entering repeated attention and feed-forward blocks, each wrapped with residual connections and normalization, and emerging as contextualized representations. Because all positions can be processed in parallel rather than sequentially, the architecture trades recurrent state for large matrix operations that map well to accelerator training.




Figure 6.9: Transformer Architecture (Encoder-Decoder): Complete architecture. The encoder (left, repeated N_L times) consists of multi-head attention followed by feed-forward layers, each with residual connections (arrows bypassing blocks) and layer normalization. The decoder (right) adds masked attention to prevent attending to future tokens during autoregressive generation. Positional encodings (sine waves) inject sequence order information absent from the permutation-invariant attention operation. This design enables training parallelism across all positions while the decoder maintains autoregressive causality during inference. Source: (Vaswani et al. 2017).

6.6.3 Computational mapping

Despite these computational costs, the effectiveness of attention has driven sustained engineering effort to push context limits ever further. Figure 6.10 is a point-in-time schematic: early widely used transformer reports exposed context windows around 512–2K tokens for BERT, GPT-2, and GPT-3-style models (Devlin et al. 2019; Radford et al. 2019; Brown et al. 2020), while later product and technical announcements reported much larger windows such as GPT-4 Turbo (128K), Claude 2.1 (200K), and Gemini 1.5 (1M+) (OpenAI 2023c; Anthropic 2023; Google 2024a). Treat these product names as dated scale anchors: the durable systems lesson is that longer contexts trade external retrieval and chunking complexity for larger attention and KV-cache budgets. Techniques like FlashAttention (T. Dao et al. 2022), sparse attention, and architectural innovations make that trade-off less expensive but do not repeal the memory scaling.




Figure 6.10: Context Window Explosion: Model context capacity (tokens, log scale) from early models (512 to 2K tokens) to modern long-context architectures (128K to 1M tokens). As context scales 1,000×, naive attention memory grows 1,000,000× ($\mathcal{O}(S^2)$), driving the development of FlashAttention, sparse attention, and linear attention mechanisms.

Listing 6.8 presents a typical implementation, showing how self-attention derives queries, keys, and values from the same input sequence.

The preceding computational mapping shows how transformers process entire sequences in parallel. The picture changes at inference time, when the model generates tokens one at a time, a shift whose system consequences section 6.6.4 quantifies through the GPT-2 XL lighthouse.

6.6.4 System implications

The quadratic bottleneck analyzed in section 6.5.4 manifests differently across execution regimes. Full-sequence training and prefill compute all token positions together, whereas autoregressive decoding generates one token at a time. Their compute, memory, and data-movement bottlenecks depend on sequence length, batch size, model dimensions, implementation, and hardware.

6.6.4.1 Training: The quadratic compute wall

During training, all token positions can be processed in parallel, while dense attention’s score computation grows as $\mathcal{O}(S^2)$. For long sequences (for example, 32k tokens), materializing the $32k \times 32k$ attention matrix requires gigabytes of memory per layer across multiple heads. This cost motivates optimizations like FlashAttention, which tiles computation to avoid materializing the full matrix in HBM. The hardware memory hierarchies (HBM, SRAM, register files) that make such tiling effective are detailed in chapter 11.

6.6.4.2 Small-batch autoregressive decoding: The memory bandwidth wall

Autoregressive decoding generates one token at a time and is often *memory-bandwidth bound* at batch one or small batch sizes. To generate a token, the system performs three broad operations:

1. Read model weights, with reuse determined by batching and the memory hierarchy (for example, a 70-billion-parameter model stores 140 GB of FP16 weights; GPT-2 XL makes the cost of that first step concrete).
2. Perform matrix-vector multiplications.
3. Read/write the KV Cache.

Listing 6.8: Self-Attention and Multi-Head Attention: Self-attention (top) derives $\mathbf{Q}$, $\mathbf{K}$, and $\mathbf{V}$ from the same input $\mathbf{X}$ through three projections, then computes attention as before. Multi-head attention (bottom) runs N_{heads} parallel attention heads with dimension $d_k = d_{\text{model}}/N_{\text{heads}}$, then concatenates and projects. For GPT-2 (768-dim, 12 heads), each head operates on 64 dimensions, partitioning the projected feature dimension across heads while enabling diverse relationship patterns.

```
def self_attention_layer(X, W_Q, W_K, W_V, d_k):
    # X: input tensor (batch_size × seq_len × d_model)
    # W_Q, W_K, W_V: weight matrices (d_model × d_k)

    Q = matmul(X, W_Q)
    K = matmul(X, W_K)
    V = matmul(X, W_V)

    scores = matmul(Q, K.transpose(-2, -1)) / sqrt(d_k)
    attention_weights = softmax(scores, dim=-1)
    output = matmul(attention_weights, V)

    return output

def multi_head_attention(X, W_Q, W_K, W_V, W_O, num_heads, d_k):
    outputs = []
    for i in range(num_heads):
        head_output = self_attention_layer(
            X, W_Q[i], W_K[i], W_V[i], d_k
        )
        outputs.append(head_output)

    concat_output = torch.cat(outputs, dim=-1)
    final_output = matmul(concat_output, W_O)

    return final_output
```

21 | KV cache memory scaling: For a 7-billion-parameter transformer in FP16, model weights consume ~14 GB; a single request's cache requires $32 \text{ layers} \times 2 \text{ (K,V)} \times 32 \text{ heads} \times 2,048 \text{ positions} \times 128 \text{ dimensions} \times 2 \text{ bytes}$, or about 1.07 GB. At 8 concurrent users, the cache is ~8.6 GB, a substantial addition to the weights, and grows linearly with context and concurrency, scaling throughput therefore forces choices such as grouped-query attention (Ainslie et al. 2023), shorter contexts, or KV paging and offloading (Kwon et al. 2023). Paging and offloading preserve model outputs, whereas grouped-query attention changes the architecture and shorter contexts change the information available to the model.

The **KV cache**²¹ grows linearly with sequence length ($\mathcal{O}(N_L \times 2 \times N_{\text{heads}} \times S \times d_{\text{head}})$ per request, distinct from the $\mathcal{O}(S^2)$ attention score matrix during training), storing the key and value vectors for all previous tokens to avoid recomputing them (Pope et al. 2023; Kwon et al. 2023). For long contexts, this cache can become massive (for example, 100+ GB), and each decoding step reads prior keys and values. The resulting bandwidth pressure depends on context length, batching, attention architecture, cache layout, and hardware.

This implementation reveals three key computational characteristics. Self-attention enables parallel processing across all positions in a full sequence. Dense score computation grows quadratically with sequence length, creating a long-sequence bottleneck. Autoregressive decoding adds a sequential dependency and can become bandwidth bound at small batch sizes.

This examination of MLPs, CNNs, RNNs, attention mechanisms, and transformers shows how each family addresses a different pattern through dense feature interactions, spatial locality, sequential dependencies, or input-dependent relationships. These families do not cover every production workload. Recommendation systems combine sparse embedding lookups, whose capacity and access distribution shape placement and caching, with dense layers that still require compute and bandwidth.

This hybrid matters at production scale: Meta reported that recommendation models accounted for most of its AI inference cycles (Gupta et al. 2020), even though such workloads receive less academic attention than language or vision models. Section 6.7 examines recommendation as the chapter's final paradigm.

 Lighthouse 6.5: GPT-2 XL (bandwidth lighthouse)

Why it matters: GPT-2 XL exemplifies **memory-bandwidth-bound** batch-one decoding. In the weight-only model used here, each decoding step reads 6 GB of FP32 weights and performs matrix-vector operations. The resulting arithmetic intensity is about 0.5 FLOP/byte with FP32 weights or 1 FLOP/byte with FP16 weights. Batching can amortize this weight traffic, while KV-cache and activation traffic add to it. Table 6.6 summarizes the bandwidth lighthouse’s quantitative properties:

Table 6.6: GPT-2 XL Systems Profile: Weight-only estimates for batch-one autoregressive decoding. Under these assumptions, each step reads the model weights at roughly 0.5–1 FLOP/byte, so HBM bandwidth strongly constrains generation throughput.

Property	Value	System Implication
Parameters	1.5B	Weight loading dominates inference latency.
Model Size	6 GB (FP32)	Fits on one GPU but saturates HBM bandwidth.
Compute	3 GFLOP/token	Low per-token compute; bottleneck is data movement, not math.
Constraint	Memory Bandwidth	Weight-bound batch-one tokens/s depends strongly on HBM bandwidth.
Profile	Bandwidth-Bound (small-batch decoding)	Larger batches can amortize weight traffic and shift the bottleneck.

6.7 Sparse Architectures: RecSys

When a user opens a streaming service, the system must select a handful of recommendations from a catalog of millions—in under 50 milliseconds. The fundamental challenge is representing both users and items as dense vectors in a shared embedding space, then computing similarity at scale.

Unlike the architectures examined so far, which are typically compute bound or bandwidth bound, recommendation models are uniquely memory-capacity-bound due to their reliance on massive embedding tables. This distinction explains why the same GPU that processes transformers efficiently may struggle with recommendation workloads.

6.7.1 Pattern processing needs

The core challenge in RecSys is handling high-cardinality categorical features. A model might need to process User IDs (billions of unique users) and Item IDs (millions of videos or products). Raw IDs carry no geometry: user 1042 is not “near” user 1043 in any meaningful sense, and a neural network cannot infer similarity from the integer itself.

Embeddings solve the representation problem by mapping each ID to a dense vector called an **embedding**²² (Mikolov et al. 2013). The systems cost is that every lookup becomes a random read into a potentially enormous table. Recommendation workloads therefore inherit a capacity and bandwidth problem before the dense neural network layers begin.

6.7.2 Algorithmic structure

The DLRM architecture (Naumov et al. 2019) standardizes this pattern as a four-stage pipeline split across dense and sparse regimes. Continuous features such as user age or time of day first flow through a bottom MLP, a compute-intensive but memory-light stage. The sparse stage then looks up categorical IDs in embedding tables, which is where the capacity wall appears:

Categorical features such as user ID or item ID are looked up in massive embedding tables. A table for one billion users with 128-dimensional vectors requires $10^9 \times 128 \times 4 \text{ bytes} \approx 512 \text{ GB}$ of memory, making this stage memory-intensive but compute-light because each lookup is essentially a memory copy. The interaction layer then combines dense vectors from the MLP with sparse embedding vectors, typically through dot products that capture user-item relationships. A top MLP finally processes the combined features to produce a probability such as click-through rate. This combination of dense and sparse computation makes DLRM the chapter’s recommendation lighthouse; section 6.7.3 quantifies the capacity-bound profile that results.

²² | **Embedding:** From the mathematical concept of embedding one space into another: neural embeddings map discrete tokens (user IDs, words) into continuous vector spaces where semantic similarity becomes geometric proximity. Word2vec popularized neural word embeddings in 2013. For systems, embedding tables create a distinctive memory access pattern: each lookup is a random read into a potentially terabyte-scale table, producing the sparse, bandwidth-bound workload that makes DLRM fundamentally different from compute-bound architectures like ResNet.

6.7.3 Computational mapping and system implications


DLRM’s computational mapping splits into two regimes that stress different hardware subsystems. The dense MLPs are standard GEMM operations, identical to the MLP computational mapping discussed in section 6.2.4 and handled efficiently by Tensor Cores. The sparse embedding lookups, however, are qualitatively different: they are index-based memory copies (gather operations) with no arithmetic, making them entirely memory-bandwidth bound at the operation level. This is distinct from the capacity constraint described earlier: the total size of the embedding tables is what makes the model memory-capacity-bound, whereas the speed of each individual gather is what makes the lookup operations memory-bandwidth bound. Because each training sample accesses a different set of embedding rows, the access pattern is effectively random, defeating caching and prefetching strategies that benefit CNNs and MLPs.

 Lighthouse 6.6: DLRM (recommendation lighthouse)

Why it matters: DLRM exemplifies **memory-capacity-bound** workloads. Its massive embedding tables often exceed the memory of a single accelerator or server, so the system must decide where those tables live before it can optimize arithmetic throughput. The interaction layer then gathers selected embedding vectors and combines them with dense features, making capacity and irregular data movement the dominant constraints. This contrasts sharply with CNNs (compute bound) and transformers (memory-bandwidth bound), requiring different hardware and deployment choices. Table 6.7 summarizes the recommendation lighthouse’s quantitative properties:

Table 6.7: DLRM Systems Profile: Quantitative properties of the recommendation lighthouse and their system consequences. Billion-entry embedding tables push the model past single-device memory, creating a capacity-bound regime that the iron law does not directly address.

Property	Value	System Implication
Embedding Parameters	25B	Parameters × 4 bytes; dominates total model size.
Model Size	100 GB (FP32)	May exceed one device’s fast memory.
Constraint	Memory Capacity	Model size > Single GPU Memory.
Bottleneck	Irregular Data Movement	Gather operations dominate sparse feature processing.
Profile	Mixed (Sparse/Dense)	Combines memory-heavy lookups with compute-heavy MLPs.




Once embedding tables exceed one device, capacity—not arithmetic—becomes the first systems constraint. A ResNet-50 (102.4 MB) or even GPT-3 (350 GB) might fit on a single node, but industrial recommendation models can reach terabytes or petabytes because of their massive embedding tables. In iron law terms (section 1.7), neither O nor D_{vol} is the binding constraint—it is raw memory capacity that limits the system, a regime the iron law was not designed to capture.

The architecture therefore breaks the single-device assumption that worked for the earlier families. A designer has three broad options, and each changes a different part of the system:

- **Shrink the tables:** Compression, hashing, or pruning can reduce capacity pressure, but these techniques may lose information about rare users or items.
- **Move the tables:** Embeddings can live in CPU memory, host memory, or a storage-backed feature system, but every lookup then pays a data-movement cost.
- **Partition the tables:** Tables can be split across multiple memory resources so no one device stores the whole model, but each request may need vectors from several partitions.

Architecturally, DLRM turns recommendation into a capacity-management problem before it becomes a compute-optimization problem (for example, at 100 million IDs and 128 FP32 values per ID, one embedding table already consumes most of an 80 GB accelerator budget). The corresponding execution strategies include training-time partitioning in chapter 8 and hardware support for fast data movement in chapter 11.



Napkin Math 6.2: The capacity wall

Problem: Consider a recommendation system for a store with 100M items using an embedding size of 128. How much memory does the item table alone require?

Math:

1. **Table entries:** 100M items.
2. **Vector size:** 128 elements.
3. **Precision:** FP32 (4 bytes per element).
4. **Table size:** $100\text{M items} \times 128 \times 4 \text{ bytes} \approx 51.2 \text{ GB}$.
5. **Capacity share:** $51.2 \text{ GB} \div 80 \text{ GB} = 64 \text{ percent}$ of device capacity.
6. **Two-table check:** Two tables at 64 percent each exceed 100 percent of device capacity.

Systems insight: A single embedding table for one feature (Items) already consumes 64 percent of an 80 GB A100 GPU. Adding a User table of the same size means the embedding tables no longer fit on a single 80 GB, motivating either smaller tables, off-device memory, or partitioning across memory resources. DLRM is *capacity bound*: our first question is not how many FLOPs the accelerator can deliver, but where the embedding state can physically reside.

Once the tables no longer fit on one device, sparse lookups cross memory and network boundaries. A request may need rows owned by several devices, so its dense MLP waits while the system gathers them. Within a node, this exchange stresses the accelerator interconnect; across nodes, it becomes many-to-many communication. The architectural consequence is that a DLRM's memory layout determines the communication pattern that serving and training must pay, creating a fleet-scale coordination problem.

Checkpoint 6.4: DLRM and sparse scatter

Recommendation systems stress a different part of the machine than CNNs or transformers.

- ☐ **Capacity-bound:** can you explain why embedding tables push DLRM into a memory capacity regime where “FLOPs” is not the binding constraint?
- ☐ **Sparse access:** can you explain why embedding lookups behave like random memory gathers (low reuse) and therefore resist caching and prefetching?
- ☐ **Placement logic:** if a single embedding table does not fit in fast device memory, can you describe which state must move, shrink, or be partitioned?
- ☐ **Data-movement bottleneck:** can you explain why sparse embedding lookups can dominate even when the dense MLP has modest FLOP cost?

The five architecture families examined earlier (MLPs, CNNs, RNNs, transformers, and DLRM-style sparse models) appear to differ fundamentally, yet they share a striking convergence: many modern deep variants reuse a small set of primitives such as dense projections, normalization, skip connections, and gating. Every transformer block contains a feedforward MLP, and gating, invented for RNNs, reappears in mixture-of-experts routing. These building blocks are portable: they originated in one architecture family but migrated broadly because the problems they solve (gradient flow, activation stability, signal routing) recur across data types and inductive biases. For systems engineers, this portability is critical because it reveals which hardware optimizations transfer across workloads and which remain architecture-specific.

6.8 Shared Building Blocks

A transformer block reuses several ideas born elsewhere: dense projections from MLPs, residual paths from deep CNNs, normalization for activation stability, and gating-like routing in later variants. The five architecture families differ in their data assumptions, but many of their engineering problems recur, so the practical question is which building blocks and optimizations transfer. Table 6.8 shows

how these primitives accumulated as architectures grew more complex: each era inherited the tools of its predecessors while adding a mechanism for the next bottleneck.

Table 6.8: Cross-Architecture Building Blocks: Engineering primitives that originated in one architecture family and now recur across many modern designs. Each building block solves a recurring problem (gradient flow, activation stability, signal routing), though not every architecture uses every primitive.

Building Block	Born In	Problem Solved	Now Used In
GEMM	MLPs	Universal function approximation	All architectures (feedforward layers)
Parameter Sharing	CNNs	Spatial efficiency	Transformers (shared projections), RNNs (weight reuse across time)
Skip Connections	ResNets (CNNs)	Gradient flow at depth	Transformers, DenseNets, U-Nets, many modern deep networks
Normalization	CNNs (BatchNorm)	Activation stability	LayerNorm (Transformers), RMSNorm (root-mean-square normalization), GroupNorm (grouped channels)
Gating	LSTMs (RNNs)	Selective signal routing	Transformers (mixture-of-experts routing), GRUs, highway networks

Those portable building blocks were shaped by the hardware available at the time. LeNet-5 (LeCun et al. 1998) trained on CPUs with networks small enough to fit in megabytes of memory. AlexNet trained its 60-million-parameter network on two GTX 580 GPUs, mapping parallel convolutions to graphics hardware (Krizhevsky et al. 2012). ResNet-152 (He et al. 2016a) became trainable because residual connections and batch normalization improved optimization at depth, using available GPU training infrastructure rather than a specific memory-capacity threshold. Transformers (Vaswani et al. 2017) became practical on contemporary GPUs and later scaled dramatically as GPU/TPU memory bandwidth and distributed training infrastructure improved. This pattern continues: each building block exploits newly available computational resources while pushing against the limits of existing systems.

6.8.1 Dense operations: The universal baseline

GEMM is the one primitive shared by every architecture in this chapter. While section 6.2 examined MLPs as dense pattern processors, the systems engineering legacy of GEMM extends far beyond MLPs. It is the feedforward layer inside every transformer block, the 1×1 pointwise convolution in MobileNets, the input and recurrent projections inside every RNN cell, and the bottom MLP in every DLRM.

MLPs introduced the GEMM-dominated computation profile that led GPU vendors to develop Tensor Cores. The backpropagation algorithm’s²³ memory access patterns, with its alternating forward and backward passes storing intermediate activations, influenced accelerator memory hierarchies. The batch processing paradigm pioneered for MLP training established the data-center-scale throughput optimization that defines modern ML infrastructure. These foundational patterns (dense matrix operations, gradient-based optimization, batch-oriented processing) appear in every architecture examined in this chapter, even when obscured by domain-specific terminology.

Dense connectivity also established the cost baseline that every subsequent architecture navigates. At $\mathcal{O}(n^2)$ parameters and operations for layers of width n , GEMM sets the reference point against which specialized architectures demonstrate efficiency gains. CNNs achieve spatial processing with $\mathcal{O}(k^2)$ parameters per location (where k is kernel size), transformers trade parameter efficiency for dynamic computation with $\mathcal{O}(S^2)$ attention complexity, and sparse architectures like DLRM exploit embedding lookups to handle categorical dimensions that would explode dense layer sizes. Each innovation represents a different strategy for escaping the dense connectivity baseline, but none escapes GEMM itself—it reappears inside every architecture as the workhorse of feature transformation.

6.8.2 Skip connections: Solving the depth problem

Parameter sharing (born in CNNs) made deep networks efficient, but efficiency alone could not solve the challenges of training them. As practitioners attempted to build deeper CNNs for more

23 | **Backpropagation:** Rumelhart, Hinton, and Williams showed in 1986 how to efficiently apply the chain rule to train multi-layer networks; standard reverse-mode training retains the forward activations needed by the backward pass, so activation memory commonly grows with network depth. Checkpointing, recomputation, and offloading can reduce stored activations by trading additional computation or data movement. Activation storage, rather than weight storage, can therefore be the binding memory constraint that determines maximum feasible batch size on a given accelerator.

complex tasks, they encountered a barrier that now confronts every deep architecture: the gradient flow problem. The mathematical foundation for skip connections starts with the failure modes of depth: vanishing gradients, exploding gradients, the limitations of ReLU, and the residual solution that enabled networks exceeding 100 layers.

6.8.2.1 The problem of depth

Backpropagation through N_L layers applies the chain rule repeatedly; section 5.4.4 gives the formal derivation of backpropagation and the chain rule. For a deep network with layers $f_1, f_2, \dots, f_{N_L}$, the gradient of the loss $\mathcal{L}$ with respect to the weights in layer 1 is:

$$\frac{\partial \mathcal{L}}{\partial W_1} = \frac{\partial \mathcal{L}}{\partial a_{N_L}} \cdot \frac{\partial a_{N_L}}{\partial z_{N_L}} \cdot \frac{\partial z_{N_L}}{\partial a_{N_L-1}} \cdot \dots \cdot \frac{\partial z_2}{\partial a_1} \cdot \frac{\partial a_1}{\partial z_1} \cdot \frac{\partial z_1}{\partial W_1}$$

where z_ℓ represents the preactivation and $a_\ell = \sigma(z_\ell)$ the postactivation output of layer ℓ . The gradient becomes a product of N_L terms, each depending on the activation function derivative $\sigma'(z_\ell)$.

Vanishing gradients create a silent training failure in deep architectures. For sigmoid activation functions, the derivative is $\sigma'(z) = \sigma(z)(1 - \sigma(z))$, with maximum value $\sigma'(0) = 0.25$. Through N_L layers, the gradient magnitude is multiplied by approximately $(0.25)^{N_L}$. With such extreme attenuation, early layers receive infinitesimal gradient signals. Weight updates become negligible, effectively preventing these layers from training.

Exploding gradients are the catastrophic counterpart to vanishing gradients. Large singular values in layer Jacobians can amplify gradient directions repeatedly through depth. For example, if each Jacobian expands a shared direction by about 1.5, that component grows exponentially. Such growth can produce numerical overflow, not a number (NaN) values, extreme parameter updates, and training divergence. Unlike vanishing gradients, which silently prevent learning, exploding gradients can cause immediate training failure.

6.8.2.2 Quantitative analysis: Plain deep networks

Consider training a deep plain convolutional network on CIFAR-10 without architectural interventions. Even with ReLU activations, which have derivative one for positive inputs, optimization can degrade as depth increases. The original ResNet paper reported that a 56-layer plain network had substantially worse CIFAR-10 test error than a 20-layer plain network (about 13.6 percent vs. 8.8 percent), demonstrating that simply adding layers can make optimization worse despite greater representational capacity (He et al. 2016a).

This “degradation problem” is not overfitting. Deeper networks train worse than shallow ones, contradicting the intuition that more layers should provide more representational capacity.

6.8.2.3 Why ReLU helps but is not sufficient

ReLU activation ($\text{ReLU}(z) = \max(0, z)$) has derivative:

$$\text{ReLU}'(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases}$$

Through active paths ($z > 0$), the derivative equals 1, avoiding gradient decay from the activation function. This represents significant improvement over sigmoid, enabling training of networks with 10–20 layers.

However, ReLU introduces a different problem: dead neurons. When $z \leq 0$, the gradient is exactly zero, blocking gradient flow through that activation. A poorly initialized neuron or large update can keep a ReLU unit negative across the training data, causing it to “die.” ReLU also does not solve gradient-flow issues arising from weight matrices themselves. Singular values far below or above one can still attenuate or amplify gradients through depth.

6.8.2.4 The residual solution

ResNet blocks introduce residual learning through skip connections that transform gradient flow. Equation 6.9 adds the identity path $\mathbf{x}$ to the residual mapping $\mathcal{F}(\mathbf{x})$.

$$\mathbf{y} = \mathcal{F}(\mathbf{x}) + \mathbf{x} \quad (6.9)$$

where $\mathcal{F}(\mathbf{x})$ represents the residual function (typically two convolutional layers with batch normalization and ReLU) and $\mathbf{x}$ is the identity skip connection.



Theorem 6.2: Residual Jacobian conditioning

Analysis: Consider the network as a composition of functions $\mathbf{x}_{\ell+1} = f_\ell(\mathbf{x}_\ell)$. By the chain rule, the gradient at the input is the product of layer Jacobians $\mathbf{J}_\ell = \frac{\partial f_\ell}{\partial \mathbf{x}_\ell}$:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{x}_0} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}_{N_L}} \cdot \prod_{\ell=1}^{N_L} \mathbf{J}_\ell$$

For plain networks, $\mathbf{J}_\ell$ is arbitrary. The bound $\|\prod_{\ell} \mathbf{J}_\ell\|_2 \leq \prod_{\ell} \|\mathbf{J}_\ell\|_2$ means uniformly subunit norms force vanishing; conversely, $\sigma_{\min}(\prod_{\ell} \mathbf{J}_\ell) \geq \prod_{\ell} \sigma_{\min}(\mathbf{J}_\ell)$ means uniformly superunit minimum singular values force expansion. Mixed, non-normal Jacobians depend on singular vectors as well as eigenvalues, so spectral radius alone does not determine gradient behavior. For ResNets, the layer function is $\mathbf{x}_{\ell+1} = \mathbf{x}_\ell + \mathcal{F}(\mathbf{x}_\ell)$, so the Jacobian is:

$$\mathbf{J}_\ell = \mathbf{I} + \frac{\partial \mathcal{F}}{\partial \mathbf{x}_\ell}$$

where $\mathbf{I}$ is the identity matrix. If $\varepsilon_\ell = \|\mathcal{F}'_\ell\|_2 < 1$, then $1 - \varepsilon_\ell \leq \sigma_{\min}(\mathbf{J}_\ell) \leq \sigma_{\max}(\mathbf{J}_\ell) \leq 1 + \varepsilon_\ell$. Thus a block whose residual Jacobian is small is close to identity and locally well conditioned. Across many blocks these deviations may still compound, so the skip connection improves gradient flow without guaranteeing unit gain or preventing cancellation.

During backpropagation, the gradient flows through this addition:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{x}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \cdot \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \cdot \frac{\partial (\mathcal{F}(\mathbf{x}) + \mathbf{x})}{\partial \mathbf{x}}$$

Applying the chain rule:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{x}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \cdot \left(\frac{\partial \mathcal{F}(\mathbf{x})}{\partial \mathbf{x}} + \mathbf{I} \right) = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \cdot \mathcal{F}'(\mathbf{x}) + \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$$

This equation reveals the critical insight. The gradient divides into two contributions. The residual contribution, $\frac{\partial \mathcal{L}}{\partial \mathbf{y}} \cdot \mathcal{F}'(\mathbf{x})$, can be small, while the identity term contributes $\frac{\partial \mathcal{L}}{\partial \mathbf{y}}$ directly. The residual contribution can still partially cancel the identity term, but when $\mathcal{F}'(\mathbf{x})$ is small, the block Jacobian remains close to identity and improves conditioning. The additive path does not guarantee unattenuated gradients through arbitrary depth.

6.8.2.5 Gradient flow through multiple blocks

Through N_L residual blocks, the gradient becomes:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{x}_0} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}_{N_L}} \cdot \prod_{\ell=1}^{N_L} (\mathcal{F}'_\ell(\mathbf{x}_\ell) + \mathbf{I})$$

Each factor $(\mathcal{F}'_\ell + \mathbf{I})$ contains the identity term, keeping the factor near identity when the residual-branch Jacobian is small. Unlike plain networks, which multiply arbitrary layer Jacobians, ResNets multiply these near-identity factors. The deviations can still compound, but this structure improves conditioning and helps make networks with 100+ layers trainable.

6.8.2.6 Empirical validation at 56 layers

The ResNet CIFAR-10 experiments provide the empirical contrast: residual networks avoided the degradation seen in deeper plain networks and achieved lower training and test error as depth increased. In the same family of experiments, a 56-layer residual network reached about 7.0 percent test error, improving over the deeper plain network rather than degrading with depth (He et al. 2016a). The critical difference appears in gradient flow and optimization: identity shortcuts give later layers a direct path to refine earlier representations rather than forcing every layer to learn a complete transformation from scratch.

Skip connections improve gradient flow but can introduce system-level costs. The addition $y = \mathcal{F}(x) + x$ requires the residual input to remain available until the paths merge. The incremental memory and runtime costs depend on tensor liveness, checkpointing, fusion, implementation, and architecture; the addition itself is usually small relative to the residual function $\mathcal{F}(x)$.

Residual networks avoided the degradation observed in comparable plain networks (He et al. 2016a). There is no fixed depth at which skip connections become necessary or insufficient: trainability also depends on initialization, normalization, optimization, layer type, and residual design (He et al. 2016b).

The gradient flow improvements from skip connections addressed one critical training challenge but left another: controlling activation distributions across layers. Even with skip connections improving gradient flow, poorly conditioned activations can destabilize training. Skip connections provide a direct gradient path; normalization can help control activation scale. This distinction explains why many deep architectures combine normalization with skip connections.

6.8.3 Normalization: Stabilizing activations at depth

Skip connections improve gradient flow to early layers; normalization helps control activation scale. Like skip connections, normalization is a portable building block: it was born as batch normalization in CNNs²⁴ (Ioffe and Szegedy 2015), evolved into layer normalization for transformers, and most recently simplified into RMSNorm²⁵ for efficient large language models. Many modern deep architectures use some variant. Understanding the mathematics of normalization reveals why these layers can materially improve deep-network training.

6.8.3.1 Batch normalization: Definition and formulation

Batch normalization normalizes activations using statistics computed over the mini-batch for each feature or channel during training. For fully connected activations, the averaging axis is the batch. For convolutional activations, implementations typically compute per-channel statistics over both the batch and spatial positions. For a mini-batch $\mathcal{B} = \{x_1, \dots, x_B\}$ of activations at a particular layer, the transformation proceeds in two stages.

First, compute the batch statistics:

$$\mu_{\mathcal{B}} = \frac{1}{B} \sum_{i=1}^B x_i \quad \sigma_{\mathcal{B}}^2 = \frac{1}{B} \sum_{i=1}^B (x_i - \mu_{\mathcal{B}})^2$$

Then, over the batch, normalize and apply learnable scale and shift. The normalization step in equation 6.10 centers and scales activations, while equation 6.11 applies learnable parameters that allow the network to recover the identity transformation if optimal:

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad (6.10)$$

$$y_i = \gamma \hat{x}_i + \beta \quad (6.11)$$

The parameters γ (scale) and β (shift) are learned during training, while ϵ (typically 10^{-5}) prevents division by zero. The affine transform restores learned scale and shift flexibility, but fixed γ and β cannot exactly invert varying statistics for every batch. Normalization often permits larger learning rates in some architectures and training regimes, which can accelerate convergence relative to an otherwise comparable unnormalized network.

²⁴ | **Batch normalization (BatchNorm)**: The original normalization layer (Ioffe and Szegedy 2015), which re-scales activations using per-mini-batch statistics. In one ImageNet experiment, it reached the target accuracy with $14\times$ fewer training steps. Its batch-size dependency and training-serving skew (switching from batch statistics to running averages at inference) are system limitations that motivated alternatives: LayerNorm removed batch dependency for transformers, and RMSNorm removed mean centering.

²⁵ | **RMSNorm (root mean square normalization)**: Introduced by Zhang and Sennrich (2019) at NeurIPS, RMSNorm simplifies LayerNorm by normalizing with the root mean square alone, dropping the mean-centering step. Across the paper's tested models and implementations, RMSNorm reduced overall running time by 7–64 percent relative to LayerNorm. LLaMA-family and Mixtral-style transformer reports use RMSNorm (Touvron, Lavril, et al. 2023; Touvron, Martin, et al. 2023; Jiang et al. 2024), illustrating why one reduction pass can matter for transformer inference latency.



Theorem 6.3: Normalization Jacobian conditioning

Analysis: The mathematical insight into why normalization aids training lies in how it conditions the Jacobian matrix of the layer. For the current batch, consider the gradient of the normalized output with respect to the input:

$$\frac{\partial \hat{x}_i}{\partial x_j} = \frac{1}{\sqrt{\sigma_B^2 + \epsilon}} \left(\delta_{ij} - \frac{1}{B} - \frac{(x_i - \mu_B)(x_j - \mu_B)}{B(\sigma_B^2 + \epsilon)} \right)$$

where δ_{ij} is the Kronecker delta. This Jacobian couples examples within the batch, with scale set by $1/\sqrt{\sigma_B^2 + \epsilon}$; the subsequent affine scale γ multiplies it. It has zero directions associated with batchwise shift invariance, so there is no universal positive eigenvalue bound such as $[0.5, 2.0]$.

Batch normalization can improve optimization empirically by standardizing intermediate scales, but it does not by itself prevent vanishing or exploding gradients through an entire network. Its quantitative effect depends on architecture, batch statistics, optimization, and parameterization rather than a universal two-to-fourfold gradient range.

6.8.3.2 Layer normalization: Architecture independence

While batch normalization enabled training of much deeper CNNs, it introduced a problematic dependency on batch statistics. This creates issues for small batch sizes (noisy statistics), varying sequence lengths (incompatible batch dimensions), and inference (requires running mean/variance estimation). Layer normalization addresses these limitations by normalizing across features rather than across the batch (Ba et al. 2016).

For an input vector $\mathbf{x} \in \mathbb{R}^{d_{\text{model}}}$ with d_{model} features:

$$\mu_{\text{LN}} = \frac{1}{d_{\text{model}}} \sum_{i=1}^{d_{\text{model}}} x_i \quad \sigma_{\text{LN}}^2 = \frac{1}{d_{\text{model}}} \sum_{i=1}^{d_{\text{model}}} (x_i - \mu_{\text{LN}})^2$$

Equation 6.12 defines the complete layer normalization operation, where $\odot$ denotes element-wise multiplication:

$$\text{LayerNorm}(\mathbf{x}) = \frac{\mathbf{x} - \mu_{\text{LN}}}{\sqrt{\sigma_{\text{LN}}^2 + \epsilon}} \odot \gamma + \beta \quad (6.12)$$

Layer normalization normalizes each sample independently, making the operation invariant to batch size and suitable for autoregressive models where per-sample independence is required (batch statistics would leak information across samples). This architectural difference explains why transformers universally adopt layer normalization: the self-attention mechanism processes sequences of varying length, and autoregressive generation requires each position to be normalized independently of batch composition.

6.8.3.3 Comparative analysis: When to use each variant

The choice between normalization variants depends on computational context. Table 6.9 summarizes the key trade-offs. BatchNorm typically stores learned scale/shift parameters plus nonlearned running mean/variance buffers; LayerNorm computes per-sample statistics at runtime and typically stores learned scale/shift parameters but no running-statistic buffers.

Batch size constraints emerge because batch normalization estimates statistics from the values available for each channel. Small effective sample counts can make those statistics noisy, but the threshold depends on architecture, spatial dimensions, task, and implementation. This constraint impacts memory-limited scenarios such as high-resolution images or large models.

The computational cost of computing mean and variance adds $\mathcal{O}(B \times d_{\text{model}})$ operations per batch normalization layer for batch size B and feature dimension d_{model} . For layer normalization, the cost is $\mathcal{O}(d_{\text{model}})$ per sample. RMSNorm reduces this further by eliminating the mean computation.

Table 6.9: Normalization Variant Comparison: Different normalization techniques trade off between computational efficiency, batch size sensitivity, and architectural compatibility. RMSNorm (Zhang and Sennrich 2019), used in LLaMA-family and other efficient transformer architectures (Touvron, Lavril, et al. 2023; Touvron, Martin, et al. 2023), omits mean centering:

$$\text{RMSNorm}(\mathbf{x}) = \mathbf{x} / \sqrt{\frac{1}{d_{\text{model}}} \sum_i x_i^2 + \epsilon \cdot \gamma}.$$

Characteristic	BatchNorm	LayerNorm	RMSNorm
Normalization Axis	Batch and, for CNNs, spatial positions	Feature dimension	Feature dimension
Batch Size Dependency	High (noisy for small batches)	None	None
Typical Use Case	CNNs, vision models	Transformers, RNNs	LLaMA, efficient transformers
Computation Cost	Higher (mean + variance)	Higher (mean + variance)	Lower (RMS only)
Training/Inference	Different (running stats)	Identical	Identical

Operational differences between training vs. inference require explicit mode switching for batch normalization, which exhibits different behavior between training (batch statistics) and inference (running statistics). Incorrect mode handling is a common source of training-serving skew. Layer normalization behaves identically in both modes, simplifying deployment.

Skip connections and normalization solve depth-related problems—gradient flow and activation stability, respectively. The third portable building block, gating, solves a different problem entirely: *selectively routing information* through the network.

6.8.4 Gating: Controlling information flow

Gating mechanisms were born in RNNs, where early sequence models hit a “temporal barrier”: gradients vanished or exploded through long sequences, revealing that simple recurrence was insufficient for long-term dependencies. LSTMs²⁶ (Hochreiter and Schmidhuber 1997) and GRUs²⁷ (Cho et al. 2014) addressed this by introducing **gates**: small MLPs that learn to control the flow of information through the network, acting as differentiable valves that selectively protect, forget, or route signals.

The key insight is that gating is not an RNN-specific technique. It is a general principle of using one learned signal to modulate another learned signal. Highway Networks applied the idea to feedforward layers, letting the network decide whether to transform an input or pass it through, which made them an important precursor to skip connections. Attention uses the same principle at the sequence level: encoder-decoder attention (Bahdanau et al. 2015), originally introduced for machine translation, learns which source positions should influence each output position. In transformers, the softmax attention weights become the routing signal that controls how much each position contributes to the output, while later large-scale variants extend the same idea to explicit expert routing. The portability of gating reinforces the central theme: the building blocks that matter most are not tied to any single architecture but solve universal problems—in this case, the problem of selectively routing information through deep, complex networks.

6.8.5 Synthesis: How transformers recombine everything

The transformer recombines several building blocks discussed earlier. A full transformer block combines the residual-path idea illustrated in figure 6.11 with dense projections, normalization, and attention gating. Dense GEMM operations in MLP-style feedforward networks process features between attention layers. Residual paths wrap every sub-layer, enabling gradient flow through deep stacks. LayerNorm, evolved from the same stabilization problem that BatchNorm addressed in CNNs, stabilizes activations at each sub-layer (Ba et al. 2016). Softmax attention weights then gate how much each position contributes, while mixture-of-experts variants extend the same routing idea to explicit expert selection.

This recombination is not accidental. The transition from RNNs to transformers represents a decisive engineering shift from sequential to parallel state management. Replacing time-step dependencies with global, data-dependent routing (attention) moves sequence models from $\mathcal{O}(S)$ sequential complexity to $\mathcal{O}(1)$ sequential steps for information flow between any two positions, enabling full use of accelerator parallelism. The other building blocks, however, carried over unchanged: GEMM, skip connections, and normalization remain essential across all families.

²⁶ | **LSTM (long short-term memory):** Invented by Hochreiter and Schmidhuber in 1997, LSTMs introduced a “Constant Error Carousel,” a gated cell state that protects error signals from exponential decay during backpropagation through time. The systems cost of this solution: a standard LSTM computes input, forget, and output gates plus a candidate cell update, giving roughly four affine transformations per time step vs. one in a vanilla RNN. This compute overhead explains why transformers, which solve long-range dependencies through parallelizable attention, replaced LSTMs in many large-scale language workloads.

²⁷ | **GRU (gated recurrent unit):** Cho et al. (2014) describes a gated hidden unit for encoder-decoder translation that uses reset and update gates to control how the hidden state is updated. Relative to an LSTM’s input, forget, and output gates plus candidate cell update, this gives a simpler gated recurrence. The broader systems lesson: architectural simplification can reduce state and matrix operations when it preserves task performance, a principle that recurs in efficiency-oriented designs from MobileNet to distilled transformers.




Figure 6.11: Residual Connection Block: The skip connection implements $y = \text{ReLU}(\mathcal{F}(x) + x)$, creating an identity shortcut that bypasses the weight layers before the final nonlinearity. During backpropagation, gradients flow through both the residual path (via $\mathcal{F}'(x)$) and the identity path (via 1 before the final nonlinearity), helping gradients reach early layers even in 100+ layer networks. This architectural pattern enabled very deep residual CNNs and became central in transformer blocks and many other modern architectures.

28 | **Vision transformers (ViTs):** Google's 2020 ViT paper split 224×224 images into 16×16 patches (196 "tokens") and applied standard transformer attention. ViTs replace CNN's local convolutions with $\mathcal{O}(S^2)$ global attention over patch tokens. In the original study, large-scale pre-training improved ViT's competitiveness, illustrating how data and compute can compensate for weaker spatial inductive bias (Dosovitskiy et al. 2021).

This portability recurs in later architectures. Vision Transformers²⁸ adapt the transformer to images while maintaining all four building blocks (Dosovitskiy et al. 2021). GPT-3, for example, scales up these transformer patterns and uses alternating dense and locally banded sparse attention while still relying on the same core primitives (Brown et al. 2020). Practical implementation challenges and optimizations are explored in chapter 10.

Table 6.10 makes this synthesis concrete. Transformers retain the core GEMM operations common to all architectures but introduce content-dependent all-to-all reductions through attention, blending the broadcast operations of MLPs with the gather and reduce operations of more dynamic architectures.

Table 6.10: Primitive Utilization Across Architectures: Every architecture uses GEMM, but they differ in memory access patterns, data movement (broadcast to gather+reduce), and parallelism. Transformers combine dense tiled QKV access with attention-specific reductions, which is why they stress both the GEMM units and the reduction path at once.

Primitive Type	MLP	CNN	RNN	Transformer
Computation	Dense GEMM	Convolution	Sequential GEMM	GEMM + Attention
Memory Access	Sequential	Strided	Sequential + State	Tiled QKV streams
Data Movement	Broadcast	Sliding window	Temporal broadcast	Gather + Reduce
Parallelism	High	High	Low (time deps)	High (positions)

For systems engineers, this building-block perspective separates portable optimizations from architecture-specific ones. GEMM tiling and mixed-precision compute benefit every architecture. Skip connection memory management applies to any residual network. Normalization kernel fusion helps CNNs and transformers alike. Attention-specific optimizations remain tied to attention's memory pattern, but even those build on the same underlying GEMM and memory-access primitives. To understand why these optimizations transfer, section 6.9 lowers the shared layers to the primitives the machine executes.

6.9 Computational Primitives

Portable building blocks still lower to a smaller set of operations; those primitives determine what hardware must execute. A ResNet-50 forward pass executes billions of multiply-accumulate operations; a transformer attention layer moves gigabytes through memory hierarchies; a DLRM lookup scatters random reads across terabyte-scale tables. Despite their architectural differences, all three reduce to a small set of *computational primitives* that hardware and software must actually execute. Synthesizing the per-architecture system implications from earlier sections into a unified view reveals common optimization opportunities.

Each primitive represents an operation that cannot be decomposed further while maintaining its essential characteristics. Understanding these operations reveals where performance bottlenecks arise on specific hardware and guides the optimization strategies detailed in chapter 11.

6.9.1 Core computational primitives

The core primitive question is which execution pattern an architecture forces the system to optimize: dense tensor math, repeated local reuse, or input-dependent routing. Matrix multiplication, sliding window operations, and dynamic computation recur across the families because each preserves a distinct performance profile when lowered to hardware. They are primitive in the engineering sense: decomposing them further would erase the performance characteristics a system must optimize.

Matrix multiplication is the dense tensor-math path. Multiplying a matrix of inputs by a matrix of weights computes weighted combinations, the core operation of neural networks (recall the reference MLP layer from section 6.2.3). This path appears everywhere: MLPs use it directly for layer computations, CNNs can lower convolutions into matrix multiplications, and transformers use it extensively in their attention mechanisms. Figure 6.12 shows one row-major arrangement: each sliding-window position unfolds into a row of the transformed matrix.




Figure 6.12: im2col Transformation: Converts convolution to GEMM by unfolding image patches into a dense matrix. The two 3×3 input feature maps produce four sliding-window positions, represented here as rows of a 4×8 matrix; the two 2×2 filter kernels stack into an 8×1 column. Their product yields the four output positions. Lowering can exploit optimized matrix kernels but expands memory by an amount that depends on kernel size, stride, padding, and implementation, so its performance benefit is workload dependent.

The **im2col**²⁹ (image to column) technique is the bridge that turns sliding-window locality into the matrix-multiply path. It unfolds overlapping image patches into a dense matrix. In the row-major arrangement in figure 6.12, each sliding-window position becomes a row and the stacked filter values form a column; transposed conventions are also common. This allows the convolution operation to be expressed as a standard GEMM operation.

The optimization trade-off is pragmatic: **im2col** spends memory to use mature GEMM implementations such as cuBLAS, MKL, and OpenBLAS. The transformation duplicates data where windows overlap, and whether lowering outperforms a direct convolution depends on the workload, implementation, and target hardware.

Structured and unstructured sparsity are treated in section 10.3.1; hardware-aware sparse execution and algorithm-hardware co-design are treated in chapter 11.

Sliding window operations are the local-reuse path. They compute local relationships by applying the same operation to chunks of data. A 3×3 convolution filter slides across the input, generating one output per window position (for example, 26×26 windows for a 28×28 input with stride 1). Modern hardware accelerators implement this through specialized memory access patterns and data buffering schemes that optimize data reuse. For example, TPUs use systolic arrays³⁰ where data flows systematically through processing elements, allowing each input value to be reused across multiple computations without repeatedly accessing off-chip memory.

Content-dependent weighting is an adaptive-routing path. In dense transformer attention, each query’s weights depend on the input, but the tensor shapes and dense computation graph remain regular and are commonly implemented with GEMMs and tiled reductions. Sparse attention and mixture-of-experts routing can additionally make the executed operations data dependent.

²⁹ | **im2col (image to column):** Rather than being a new learning algorithm, **im2col**-style lowering is an implementation technique used by CNN frameworks and libraries such as Caffe and cuDNN (Jia et al. 2014; Chetlur et al. 2014): it converts convolutions into standard GEMM calls by unfolding overlapping patches into matrix columns. The trade-off is memory: in a simple fully materialized stride-1 $K \times K$ transform, interior input elements can appear in up to K^2 columns (9 times for 3×3 filters), though borders, stride, padding, tiling, and direct-convolution algorithms reduce the realized expansion. This memory-for-simplicity exchange explains why mobile frameworks (TFLite, NNAPI) prefer direct convolution, while data center GPUs with abundant HBM may use GEMM-oriented lowering when it improves throughput.

³⁰ | **Systolic array:** Named for the heart’s rhythmic contraction, the array’s lockstep “pulse” of data through a grid of processors directly implements the efficient data reuse required by sliding window operations. By passing input values between neighboring processors, an expensive round-trip to off-chip DRAM is avoided for every single multiplication in the convolution. This is critical for efficiency, as a single off-chip memory access can cost over $100\times$ more energy than a floating-point multiply-accumulate, and still more relative to low-precision arithmetic.

Real architectures combine these paths, which is why primitive-level reasoning is a design tool rather than a taxonomy. A transformer layer uses matrix multiplications of shape $[S, d_{\text{model}}] \times [d_{\text{model}}, d_{\text{proj}}]$ for feature projections and computes an $S \times S$ score pattern for dense attention. Some variants use sliding windows or data-dependent sparse routing. The interaction between primitives creates specific demands on system design, from memory hierarchy organization to computation scheduling.

The core computational primitives above explain why certain hardware features exist (Tensor Cores for matrix multiplication) and why software frameworks organize computations in particular ways (batching similar operations together). Computational primitives, however, tell only part of the story: the way operations access memory often determines real-world performance more than the operations themselves.

6.9.2 Memory access primitives

The next optimization decision is whether the primitive can feed the compute units predictably. Memory access often constitutes the primary bottleneck in ML systems: even a matrix multiplication unit capable of thousands of operations per cycle will remain idle if data is not available in time. Accessing data from DRAM typically requires hundreds of cycles, while on-chip computation requires only a few, making data movement a first-order energy constraint.

The relevant access patterns are sequential access, strided access, and random access because they determine how much of the memory stream the system can predict and reuse. Each pattern creates different demands on the memory system and offers different opportunities for optimization. Critically, each incurs vastly different energy costs based on the preceding principle.

Systems Perspective 6.2: The energy cost of data movement

A core systems principle: *moving data costs more than computing on it*. A floating-point multiply-accumulate operation consumes approximately 1–5 pJ depending on process node (for example, ~4.6 pJ at 45 nm (Horowitz 2014)), while fetching a 32-bit operand from off-chip DRAM costs hundreds of pJ in the same reference model. Even conservative estimates put off-chip DRAM access tens to hundreds of times above arithmetic, explaining why memory access patterns, not just FLOP counts, determine real-world efficiency.

Revisit the preceding architectures through this energy lens: MLPs have low data reuse (each weight loaded once per sample) and are therefore energy-dominated by DRAM traffic. CNNs reuse filter weights across spatial positions, amortizing load cost over $H \times W$ applications; the very locality that makes them compute-bound also makes them energy-efficient. RNNs reuse weights across time steps (high temporal reuse) but pay repeated hidden-state read/write costs at each step. Transformers combine pairwise score computation with key-value movement, making dense full-sequence attention compute-quadratic; implementation and tiling determine auxiliary memory traffic and energy. These energy profiles directly track the bottleneck column in table 6.3.

This principle underlies later optimization strategies: section 10.4 shows how quantization reduces bits moved per value, pruning eliminates unnecessary data movement, and tiling keeps working sets in faster, lower-energy caches.

Sequential access is the simplest, most efficient pattern, and the most energy-favorable. Consider an MLP performing matrix multiplication: it accesses weight matrices and input vectors in contiguous order. This pattern maps well to modern memory systems; DRAM can operate in burst mode for sequential reads (reaching on the order of hundreds of GB/s in modern GPUs), and hardware prefetchers can effectively predict and fetch upcoming data. Software frameworks optimize for this by ensuring data is laid out contiguously in memory and aligning data to cache line boundaries.

Strided access appears prominently in CNNs, where each output position needs to access a window of input values at regular intervals. Each output position requires accessing nine input values (for a 3×3 filter) with a stride matching the input width. While less efficient than sequential access, hardware supports this through pattern-aware caching strategies and specialized memory controllers. Software frameworks often transform these strided patterns into sequential access

through data layout reorganization, where the im2col transformation in deep learning frameworks converts convolution's strided access into efficient matrix multiplications.

Random access poses the greatest challenge for system efficiency. Sparse embedding lookups in recommendation models illustrate this challenge. Each request may touch a different set of table rows, defeating predictable streaming and causing cache misses or irregular memory latency. Dense transformer attention is different. Its weights are content-dependent, but Q, K, and V are stored in contiguous tensors and optimized kernels tile those tensors through SRAM/registers while reducing over the sequence. The systems challenge for dense attention is therefore quadratic score computation and reduction structure, not arbitrary address-random fetches.

Table 6.11 quantifies how these different memory access patterns contribute to the overall memory requirements of each architecture, comparing MLPs, CNNs, RNNs, and transformers across parameter storage, activation storage, and scaling behavior.

Table 6.11: Memory Access Complexity: Different neural network architectures exhibit varying memory access patterns and storage requirements, impacting system performance and scalability. Parameter storage depends on model dimensions, while activation storage represents a significant runtime cost. The transformer's $\mathcal{O}(BS^2)$ term describes materialized attention scores in a naïve implementation; tiled exact-attention kernels can avoid retaining that full matrix.

Architecture	Input Dependency	Parameter Storage	Activation Storage	Scaling Behavior
MLP	Linear	$\mathcal{O}(N_{\text{in}} \times d_{\text{width}})$	$\mathcal{O}(B \times d_{\text{width}})$	Predictable
CNN	Constant w.r.t. resolution	$\mathcal{O}(K^2 C_{\text{in}} C_{\text{out}})$	$\mathcal{O}(B \times H_{\text{img}} \times W_{\text{img}} \times C)$	Efficient
RNN	Linear	$\mathcal{O}(d_{\text{hidden}}^2)$	$\mathcal{O}(B \times S \times d_{\text{hidden}})$	Challenging
Transformer	Quadratic attention	$\mathcal{O}(d_{\text{model}}^2 + d_{\text{model}} d_{\text{ff}})$ per block	$\mathcal{O}(B \times S^2)$ attention, plus $\mathcal{O}(BS d_{\text{model}})$ activations	Problematic

Where:

- N_{in} : Input size
- d_{width} : Layer width
- B : Batch size
- K : Kernel size
- C : Number of channels
- $C_{\text{in}}, C_{\text{out}}$: Input and output channels
- H_{img} : Height of input feature map (CNN)
- W_{img} : Width of input feature map (CNN)
- d_{hidden} : RNN hidden-state dimension
- S : Sequence length
- d_{model} : Transformer model dimensionality

Table 6.11 captures *where* data lives and how access patterns scale. The complementary table 6.12 that follows captures *how much* computation each architecture demands, including forward-pass FLOPs, parallelization potential, and the resulting bottleneck. Together, the two tables answer the systems questions “how much work?” and “how does the memory system handle it?”, providing a resource profile that informs design decisions such as choosing memory hierarchy configurations and developing memory optimization strategies.

The impact of these patterns becomes clear when we consider data reuse opportunities. In CNNs, each input pixel participates in multiple convolution windows (typically nine times for a 3×3 filter), making effective data reuse necessary for performance. Modern GPUs provide multi-level cache hierarchies (L1, L2, shared memory) to capture this reuse, while software techniques like loop tiling ensure data remains in cache once loaded.

Working set size, the amount of data needed simultaneously for computation, varies dramatically across architectures. An MLP layer might need only a few hundred KB (weights plus activations), while a transformer processing long sequences can require several MB just for storing attention patterns. These differences directly influence hardware design choices, like the balance between

compute units and on-chip memory, and software optimizations like activation checkpointing, which saves memory by recomputing selected activations during backpropagation instead of storing all of them, or attention approximation techniques.

Table 6.12: Computational Complexity Comparison: Scaling behaviors and resource requirements for major neural network architectures. Variables: d_{in} = input dimension, d_{out} = output dimension, d_{model} = transformer model dimension, d_{hidden} = RNN hidden-state dimension, k = kernel size, c = channels, H_{img} , W_{img} = spatial dimensions, S = sequence length (time steps), B = batch size. The bottleneck column identifies the primary performance-limiting factor in typical deployments. For transformers, the $\mathcal{O}(S^2)$ memory term assumes materialized attention scores; tiled exact-attention kernels can avoid retaining that full matrix.

Architecture	Parameters	Forward Pass	Memory	Parallelization	Bottleneck
MLPs	$\mathcal{O}(d_{\text{in}} \times d_{\text{out}})$ per layer	$\mathcal{O}(d_{\text{in}} \times d_{\text{out}})$ per layer	$\mathcal{O}(d_{\text{in}} d_{\text{out}})$ weights $\mathcal{O}(B d_{\text{out}})$ activations	Excellent Matrix ops parallel	Memory bandwidth
CNNs	$\mathcal{O}(k^2 \times c_{\text{in}} \times c_{\text{out}})$ per layer	$\mathcal{O}(H_{\text{img}} \times W_{\text{img}} \times k^2 \times c_{\text{in}} \times c_{\text{out}})$	$\mathcal{O}(H_{\text{img}} \times W_{\text{img}} \times c)$ features $\mathcal{O}(k^2 \times c^2)$ weights	Good Spatial independence	Often compute throughput; bandwidth for depthwise or small-batch cases
RNNs	$\mathcal{O}(d_{\text{hidden}}^2 + d_{\text{hidden}} \times d_{\text{in}})$ total	$\mathcal{O}(S \times d_{\text{hidden}}^2)$ for S time steps	$\mathcal{O}(d_{\text{hidden}})$ recurrent state (inference); $\mathcal{O}(S d_{\text{hidden}})$ activations (training)	Poor Sequential deps	Sequential deps
Transformers	$\mathcal{O}(d_{\text{model}}^2)$ QKV/O projections plus $\mathcal{O}(d_{\text{model}} d_{\text{ff}})$ feed-forward layers	$\mathcal{O}(S^2 \times d_{\text{model}} + S \times d_{\text{model}}^2)$ per layer	$\mathcal{O}(S^2)$ attention $\mathcal{O}(S \times d_{\text{model}})$ sequences	Excellent (positions) Limited by memory	Memory (S^2)

Understanding these memory access patterns is essential as architectures evolve. The shift from CNNs to transformers, for instance, has driven the development of hardware with larger on-chip memories and more advanced caching strategies to handle increased working sets and more dynamic access patterns. Future architectures will likely continue to be shaped by their memory access characteristics as much as their computational requirements.

6.9.3 Data movement primitives

Memory access patterns describe where data resides, but a complementary dimension determines system performance: the flow of information between components. Data movement primitives characterize these flows. As established in section 6.9.2, data movement often dominates both time and energy budgets, making these flow patterns critical optimization targets.

The data-movement decision is fan-out and fan-in: whether one value must reach many consumers, many values must converge, or different values must be routed to different destinations. Figure 6.13 separates the four recurring patterns: broadcast, scatter, gather, and reduction. Broadcast operations send the same data to multiple destinations simultaneously. In matrix multiplication with batch size 32, each weight must be broadcast to process different inputs in parallel. Modern hardware supports this through specialized interconnects and hardware multicast capabilities, with bandwidth on the order of hundreds of GB/s in high-end accelerator interconnects, while some accelerators also use dedicated on-chip broadcast fabrics. Software frameworks optimize broadcasts by restructuring computations (like matrix tiling) to maximize data reuse.

Scatter operations distribute different elements to different destinations. When parallelizing a 512×512 matrix multiplication across accelerator cores, each core receives a subset of the computation. This parallelization is important for performance but challenging, as memory conflicts and load imbalance can reduce efficiency substantially. Hardware provides flexible high-bandwidth interconnects (often in the hundreds of GB/s class within a node), while software frameworks employ specialized work distribution algorithms to maintain high utilization. In large language models, mixture-of-experts architectures expose a far more demanding scatter pattern: a learned gating network routes each token to a small subset of expert sub-networks distributed across accelerators, requiring all-to-all communication that scales with the number of devices. Unlike the predictable tile-to-core scatter in matrix tiling, expert routing is data-dependent, so load imbalance across experts is common and can leave most accelerator capacity idle while a handful of hot experts




Figure 6.13: Data Movement Primitives: Four fundamental patterns govern information flow in neural network computation. Broadcast (top-left) replicates a single value to all destinations, used when sharing weights across batch elements. Scatter (top-right) distributes distinct elements to different destinations, enabling work partitioning. Gather (bottom-left) collects distributed values to a single location, as in attention pooling. Reduction (bottom-right) combines multiple values through aggregation (sum, max), appearing in gradient synchronization and attention scoring. The energy cost of data movement (see section 6.9.2) makes these patterns critical optimization targets.

become bottlenecks. This communication cost is a primary constraint on scaling such models to hundreds of experts.

Gather operations collect data from multiple sources. Dense transformer attention combines information from every key-value position, but implementations operate on regular dense tensors and commonly tile these reductions for locality. Irregular random gathers instead arise in workloads such as sparse embeddings, graph neighborhoods, and data-dependent routing.

Reduction operations combine multiple values into a single result through operations like summation. When computing attention scores in transformers or layer outputs in MLPs, efficient reduction is essential. Hardware implements tree-structured reduction networks (reducing latency from $\mathcal{O}(n)$ to $\mathcal{O}(\log n)$), while software frameworks use optimized parallel reduction algorithms that can achieve near-theoretical peak performance.

In practice, these patterns combine in layered ways. For each sequence and attention head in a transformer attention operation with sequence length 512 and batch size 32, the computation involves broadcasting query vectors (512×64 elements), gathering relevant keys and values ($512 \times 512 \times 64$ elements), and reducing attention scores (512×512 elements). The batch dimension multiplies each of these counts by 32.

The evolution from CNNs to transformers has increased reliance on gather and reduction operations, driving hardware innovations like more flexible interconnects and larger on-chip memories. As models grow (some now exceeding 100 billion parameters), efficient data movement becomes an architecture constraint rather than an implementation afterthought, leading to innovations like near-memory processing and targeted data flow optimizations.

6.9.4 System design impact

The computational, memory access, and data movement primitives explored earlier become system design constraints when they force resources to be allocated in silicon and software. Matrix-heavy workloads justify tensor units; random and gather-heavy workloads justify memory hierarchy and interconnect investment. The way these primitives influence hardware design, create common bottlenecks, and drive trade-offs turns architecture selection into infrastructure planning.

The most visible result is specialized hardware. The prevalence of matrix multiplications and convolutions in deep learning has led to the development of TPUs³¹ and Tensor Cores in GPUs,

³¹ | **TPU (tensor processing unit):** Google's first TPU maps matrix multiplication onto a large systolic array, trading general-purpose features such as caches and complex control flow for domain-specific inference efficiency (Jouppi et al. 2017). The architectural lesson needed here is qualitative: when one primitive dominates a workload, dedicated data paths and local reuse can outperform general-purpose flexibility. Chapter 11 develops the precision choices, hardware specifications, and performance trade-offs.

which are specifically designed to perform these operations efficiently. Chapter 11 examines how these specialized units map architectural primitives to silicon, from systolic arrays for GEMM to dataflow engines for convolution.

Memory systems have also evolved in response to deep learning primitives. The need to support both sequential and random access patterns efficiently has driven the development of multi-level memory hierarchies. **HBM**—3D-stacked DRAM delivering 2–3 TB/s of bandwidth, over 20× standard server RAM—has become common in AI accelerators to support high data-movement requirements, especially for operations such as transformer attention. On-chip memory hierarchies have grown in complexity, with multiple levels of caching and **scratchpad memories**—programmer-controlled SRAM that trades cache convenience for explicit data movement and predictable locality—to support the diverse working set sizes of different neural network layers.

The data movement primitives have particularly influenced the design of interconnects and on-chip networks. The need to support efficient broadcasts, gathers, and reductions has led to the development of more flexible and higher-bandwidth interconnects. Some AI chips now feature specialized networks-on-chip designed to accelerate common data movement patterns in neural networks.


The system implications of these primitives span hardware, software, and performance considerations. Table 6.13 turns the primitive-to-system mapping into a design checklist: each row links an architectural primitive to the hardware support, software optimization, and bottleneck it tends to create. Despite the specialized hardware that these primitives have motivated, several bottlenecks persist. Memory bandwidth often remains a key limitation, particularly for models with large working sets or those that require frequent random access. The energy cost of data movement, especially between off-chip memory and processing units, continues to be a significant concern. For large-scale models, the communication overhead in distributed training can become a bottleneck, limiting scaling efficiency.

Table 6.13: Primitive-Hardware Co-Design: Efficient machine learning systems require tight integration between algorithmic primitives and underlying hardware. Common primitives map to specific hardware accelerations and software optimizations, each presenting distinct implementation challenges. Specialized hardware such as Tensor Cores addresses computational demands of matrix multiplication and sliding windows, while software techniques like batching and dynamic graph execution further enhance performance.

Primitive	Hardware Impact	Software Optimization	Key Challenges
Matrix Multiplication	Tensor Cores	Batching, GEMM libraries	Parallelization, precision
Sliding Window	Specialized datapaths	Data layout optimization	Stride handling
Dynamic computation	Flexible routing	Dynamic graph execution	Load balancing
Sequential Access	Burst mode DRAM	Contiguous allocation	Access latency
Random Access	Large caches	Memory-aware scheduling	Cache misses
Broadcast	Specialized interconnects	Operation fusion	Bandwidth
Gather/Scatter	High-bandwidth memory	Work distribution	Load balancing

The same primitive mapping also determines energy budgets. Each architectural pattern exhibits distinct energy characteristics that inform deployment decisions and optimization strategies for data center and edge systems.

Large batched GEMMs in MLPs can achieve excellent arithmetic intensity, but small-batch MLP inference often has low reuse and can spend much of its energy on data movement. The reference FP32 multiply costs approximately 3.7 pJ/FLOP, while one 32-bit DRAM access costs 640 pJ (Horowitz 2014), about 173× as much per event. This ratio alone does not prove total dominance because workload energy also depends on operation counts, transfer counts, and reuse. Data movement can nevertheless dominate low-reuse inference, making bandwidth and locality important energy targets. This energy gap has driven accelerator designers to maximize on-chip SRAM capacity, keeping frequently reused weights and activations closer to compute and avoiding the DRAM penalty. Architectures that keep their working sets or active tiles on-chip, whether through tiled SRAM banks or wafer-scale integration, reduce inference energy by avoiding repeated off-chip transfers.



Convolutional operations reduce energy consumption through data reuse but exhibit variable efficiency depending on implementation. Im2col-based convolution implementations trade memory for simplicity; a fully materialized lowering can multiply temporary storage and memory traffic, up to K^2 for stride-1 $K \times K$ filters away from the borders. Direct convolution implementations can achieve substantially better energy efficiency by eliminating redundant data movement, particularly for larger kernel sizes where im2col duplication is most severe.

Sequential processing in RNNs creates energy efficiency opportunities through temporal data reuse. The constant memory footprint of RNN hidden states allows aggressive caching strategies that can dramatically reduce DRAM access energy for long sequences by keeping the recurrent state in on-chip SRAM. The sequential dependencies limit parallelization opportunities, often resulting in suboptimal hardware utilization and higher energy per operation.

Attention mechanisms in transformers can exhibit high energy consumption per operation due to data movement and, in naive implementations, stored attention matrices (the quadratic bottleneck from section 6.5.4). Attention's data movement can raise energy per useful FLOP compared with standard matrix multiplication, making long-sequence processing expensive without implementations such as FlashAttention.

These energy profiles make primitive support a deployment trade-off. Optimizing for the dense matrix operations common in MLPs and CNNs might come at the cost of flexibility needed for the more dynamic computations in attention mechanisms. Supporting large working sets for transformers might require sacrificing energy efficiency.

The right balance depends on the target workloads and deployment scenarios. Understanding the nature of each primitive guides the development of both hardware and software optimizations in ML systems, allowing designers to make informed decisions about system architecture and resource allocation.

The analysis of architectural patterns, computational primitives, and system implications provides the conceptual foundation for understanding *how* architectures work and *what* they cost. The practical selection problem is to choose an architecture for a specific problem under specific deployment constraints. This selection process must consider not only algorithmic performance but also the deployment constraints covered in chapter 2 and the lifecycle requirements introduced in chapter 3.

6.10 Architecture Selection Framework

A wildlife monitoring sensor may need to classify camera-trap images on solar power, while a recommendation service may need to retrieve terabyte-scale embeddings under a millisecond budget. Those deployment constraints immediately rule out many otherwise accurate architectures. The families examined earlier embody specific assumptions about data structure and computational patterns: MLPs assume arbitrary feature relationships, CNNs exploit spatial locality, RNNs capture temporal dependencies, and transformers model complex relational patterns. The selection problem is to match those assumptions to a specific use case before optimizing the model.

Successful architecture selection requires understanding principles rather than following trends: matching data characteristics to architectural strengths, evaluating computational constraints against system capabilities, and balancing accuracy requirements with deployment realities. The framework presented here draws upon the computational patterns and system implications explored in this chapter, together with the deployment paradigms from chapter 2 and the lifecycle constraints from chapter 3. The same selection logic also governs chapter 9 and chapter 14, where data curation and production operations add their own constraints.

6.10.1 Data-to-architecture mapping

The first step in systematic architecture selection involves data-to-architecture mapping: understanding how different data types align with architectural strengths. The architectural families introduced in section 6.1 provide the foundation: MLPs for tabular data with arbitrary relationships, CNNs for spatial data with local patterns, RNNs for sequential data with temporal dependencies, transformers for complex relational data where any element might influence any other, and sparse embedding architectures such as DLRM for high-cardinality categorical recommendation data.

This alignment is not coincidental; it reflects fundamental computational trade-offs. Architectures that match data characteristics can exploit natural structure for efficiency, while mismatched

architectures must work against their design assumptions, leading to poor performance or excessive resource consumption.

In practice, MLPs excel for financial modeling, medical measurements, and structured prediction where feature relationships are unknown *a priori*. CNNs dominate image recognition, 2D sensor processing, and signal analysis where spatial locality matters. RNNs remain useful for time-series forecasting and simple sequential tasks where memory across time is essential. Transformers are widely used for machine translation and language understanding (Vaswani et al. 2017; Devlin et al. 2019), and large variants support prompting-based reasoning tasks (Wei et al. 2022). DLRM-style sparse architectures are the natural starting point for recommendation systems with user IDs, item IDs, and other high-cardinality categorical features whose embedding tables dominate memory capacity.

Beyond data type matching, computational constraints often determine final feasibility. Understanding the scaling behavior of each architecture allows realistic resource planning and prevents costly architectural mismatches during deployment.

6.10.2 Computational complexity considerations

Architecture selection must account for computational and memory trade-offs that determine deployment feasibility. Each architecture exhibits distinct scaling behaviors that create different bottlenecks as problem size increases, and understanding these patterns allows realistic resource planning.

The preceding sections analyzed each architecture through the four-part lens of pattern processing needs, algorithmic structure, computational mapping, and system implications. As table 6.12 showed earlier alongside table 6.11, examining these architectures from both computational scaling and memory access perspectives reveals different optimization opportunities and system design considerations.

6.10.2.1 Scalability and production considerations

Production deployment introduces constraints beyond algorithmic performance: latency requirements, memory limitations, energy budgets, and fault tolerance needs. These are not four independent scorecards. Each family's production behavior traces back to the single structural property that defined it: dense connectivity, spatial locality, sequential dependence, or all-to-all attention. The same property that set a family's accuracy also governs how it parallelizes, how its latency scales, and how much memory it consumes.

MLPs and CNNs occupy the easier operational corner because they are largely stateless across examples and can scale when independent inputs are split across devices. Their latency and memory behavior still differ. MLP latency is usually predictable from layer size, which helps it meet strict service level agreements, while CNN latency depends more on implementation strategy, convolution algorithm, model shape, precision, and hardware support. MLPs require fixed memory proportional to model size; CNNs add feature-map memory that grows with input resolution.

RNNs and transformers create harder production regimes for opposite reasons. RNNs keep a compact hidden state, but time step t depends on time step $t - 1$, so additional hardware cannot remove the sequence's critical path and temporal state complicates recovery. Transformers parallelize well across sequence positions and deliver high throughput for batches, but the quadratic attention bottleneck (section 6.5.4) limits effective batch size, single-request latency, and checkpoint practicality as model scale grows. Hardware efficiency varies with operation shape, batch size, implementation, and target accelerator; sequential dependencies generally constrain RNN utilization, while small transformer batches often use compute units less efficiently than large batches. Chapter 8 later formalizes the corresponding scaling strategies as data, model, pipeline, and tensor parallelism.

6.10.2.2 Hardware mapping and optimization strategies

Different architectural patterns require distinct optimization strategies for efficient hardware mapping, so performance tuning starts by matching the operation shape to the hardware path. Dense matrix operations in MLPs map naturally to tensor processing units and GPU Tensor Cores (chapter 11 details how these map to specific silicon implementations). These operations benefit from three recurring optimizations: matrix tiling keeps active blocks close to the compute units, often with tile

sizes such as 64×64 for L1 cache, 256×256 for L2 cache, and 16×16 Tensor Core blocks on Volta-class GPUs; mixed-precision computation increases useful operations per second when accuracy allows it; and operation fusion reduces memory traffic by combining adjacent steps. Chapter 7 later examines how frameworks translate these high-level operations into optimized kernel launches on specific hardware.

CNNs benefit from specialized convolution algorithms and data layout optimizations that differ significantly from dense matrix operations. Im2col transformations convert convolutions to matrix multiplication but can multiply temporary storage and memory traffic, up to K^2 for fully materialized stride-1 $K \times K$ filters away from the borders. Winograd algorithms³² reduce multiplication count by $2.25\times$ for 3×3 convolutions but can amplify numerical error. Direct convolution with custom kernels can avoid im2col materialization but requires architecture-specific tuning.

RNNs require different optimization approaches because, as section 6.4.4 established, their sequential critical path cannot be shortened by adding hardware, so the available levers attack the overhead around that path instead. Loop unrolling removes per-step control overhead, shaving the latency term at the cost of larger code size and activation memory. State vectorization batches multiple independent sequences through the same step, recovering SIMD throughput without shortening any single sequence’s critical path. Wavefront parallelization exploits the independence of forward and backward passes in bidirectional models, roughly doubling utilization where the model structure permits it. None of these removes the sequential dependency; they amortize or sidestep it.

Transformer attention demands specialized optimizations that reduce memory usage and complexity. The common theme is to keep attention scores close to the compute units or to avoid computing scores that the model structure does not need. Section 8.6.4 examines FlashAttention³³ as a concrete tiling example, while sparse attention patterns remain a model-structure optimization.

The complexity patterns detailed in each architecture’s System Implications section define its most favorable domains. MLPs excel when parameter efficiency is not critical, CNNs dominate for moderate-resolution spatial data, RNNs remain viable for very long sequences where memory is constrained, and transformers excel for complex relational tasks where their computational cost is justified through superior performance. These constraints supply the inputs to a systematic decision framework for architecture selection.

6.10.3 Decision framework

Effective architecture selection requires balancing multiple competing factors: data characteristics, computational resources, performance requirements, and deployment constraints. In practice, teams often make this choice based on familiarity (“we always use transformers”) or trend-following (“new papers use X”), leading to architectures that are either overpowered for the problem (wasting resources) or underpowered (failing to meet requirements). While data patterns provide initial guidance and complexity analysis establishes feasibility bounds, final architectural choices often involve nuanced trade-offs demanding systematic evaluation.

The decision flowchart in figure 6.14 begins by identifying the data type, branches to candidate dense architectures (Transformers, RNNs, CNNs, or MLPs), and then checks each constraint diamond. High-cardinality recommendation workloads sit outside this flowchart and should be routed to the sparse embedding/DLRM family (section 6.7) before applying the same memory, compute, speed, accuracy, and deployment checks. If any check fails, the “No” path loops back for reconsideration. This iterative structure ensures consideration of all relevant factors while avoiding selection based on novelty or perceived sophistication.

When constraints require scaling down, the model compression techniques in chapter 10 provide systematic approaches for reducing memory, compute, and latency while preserving accuracy. The framework applies through three ordered steps:

1. **Data analysis:** Pattern types in data provide the strongest initial signal. Spatial data naturally aligns with CNNs, sequential data with RNNs.
2. **Progressive constraint validation:** Each constraint check (memory, computational budget, inference speed) acts as a filter. Failing any constraint requires either scaling down the current architecture or considering a fundamentally different approach.

³² | **Winograd algorithm:** For one 2×2 output tile with a 3×3 filter, this method trades 36 direct multiplications for 16 elementwise multiplications in the Winograd domain, plus additional transforms and additions. The transforms can amplify rounding error, so low-precision suitability depends on the variant, implementation, and accuracy requirements.

³³ | **FlashAttention:** An IO-aware algorithm (T. Dao et al. 2022) that avoids materializing the full $S \times S$ attention matrix in HBM by fusing computation into a single kernel tiled to fit in SRAM. The result: $2\text{--}4\times$ wall-clock speedup and memory reduction from $\mathcal{O}(S^2)$ to $\mathcal{O}(S)$, enabling training on sequences $4\text{--}16\times$ longer than standard attention. FlashAttention demonstrates that algorithmic optimization of data movement (D_{vol}) can yield larger speedups than increasing raw compute (R_{peak}) – a concrete validation of the iron law’s data term.




Figure 6.14: Architecture Selection Decision Framework: A systematic flowchart for choosing dense neural network architectures based on data characteristics and deployment constraints. The process begins with data type identification (text/sequences/images/tabular) to select initial architecture candidates (Transformers/RNNs/CNNs/MLPs), then iteratively evaluates memory budget, computational cost, inference speed, accuracy targets, and hardware compatibility. High-cardinality recommendation workloads are handled separately by the sparse-embedding family (DLRM) in section 6.7, then re-enter the same constraint checks.

- Iterative trade-off handling:** When accuracy targets remain unmet, additional model capacity may be needed, requiring a return to constraint checking. If deployment hardware cannot support the chosen architecture, reconsidering the entire architectural approach may be necessary.

6.10.4 Inductive bias hierarchy

The decision framework provides practical guidance for architecture selection, but the entire process rests on a deeper unifying principle: diverse architectures form a spectrum of structural constraints, or inductive biases. The five architectural families, practical selection framework, and computational primitives examined throughout this chapter share this common theoretical foundation, introduced

in section 6.1. Comparing their biases reveals a hierarchy of systems implications without redefining each term.

Different architectures form a hierarchy of decreasing inductive bias. CNNs exhibit the strongest constraints through local connectivity, parameter sharing, and translation equivariance, dramatically reducing the parameter space while limiting flexibility to spatial data. RNNs demonstrate moderate bias through sequential processing and shared temporal weights. MLPs maintain minimal architectural bias, requiring more data to learn structure that other architectures encode explicitly. Transformers represent adaptive inductive bias, dynamically adjusting based on data through learned attention patterns.

Many deep architectures implement hierarchical representation learning, but through different mechanisms: CNNs through progressive receptive field expansion (section 6.3), RNNs through hidden state evolution (section 6.4), and transformers through multi-head attention (section 6.5). This hierarchical organization reflects a general principle: complex patterns can be efficiently represented through composition of simpler components. For systems engineering, computational patterns must efficiently compose lower-level features into higher-level abstractions, memory hierarchies must align with representational hierarchies to minimize data movement, parallelization strategies must respect hierarchical dependency structure, and hardware accelerators must efficiently support the matrix operations implementing feature composition.

6.10.5 Architecture selection in practice

A real-time wildlife monitoring scenario synthesizes the chapter's concepts by applying the full architecture-selection process. First, a back-of-the-napkin calculation reveals the throughput ceiling that drives the hardware selection.

Napkin Math 6.3: The throughput ceiling

Problem: A real-time application must process 30 FPS of video with ResNet-50. What sustained compute throughput is required?

Math:

1. **Model cost:** ResNet-50 requires ~ 8.2 GFLOP per 224×224 image.
2. **Frame rate:** 30 FPS required.
3. **Sustained throughput:** $30 \text{ FPS} \times 8.2 \text{ GFLOP} = 246 \text{ GFLOP/s}$.
4. **Effective throughput:** $10 \text{ TFLOP/s peak} \times 50\text{--}60 \text{ percent utilization} = 5 \text{ TFLOP/s--}6 \text{ TFLOP/s}$.
5. **ResNet-50 headroom:** $5 \text{ TFLOP/s} \div 246 \text{ GFLOP/s} = 20.3\times$.
6. **Detection-model check:** $30 \text{ FPS} \times 100 \text{ GFLOP} = 3 \text{ TFLOP/s}$, leaving $1.7\times$ headroom at the lower effective-throughput estimate.

Systems insight: A mid-range GPU delivering 10 TFLOP/s theoretical peak achieves $\sim 50\text{--}60$ percent utilization in this planning scenario, yielding $5 \text{ TFLOP/s--}6 \text{ TFLOP/s}$ effective. For ResNet-50 at 30 FPS, the system has $20.3\times$ headroom. Switching to an object detection model at 100 GFLOP per frame, however, requires 3 TFLOP/s sustained, leaving only $1.7\times$ headroom. Batch size constraints or multi-stream processing quickly push the system toward the compute ceiling. ResNet-50 is compute-bound, but the margin depends on the accelerator and utilization achieved.

The throughput ceiling converts an abstract compute requirement into a concrete hardware utilization percentage. A real-world deployment adds physical constraints the ceiling alone does not capture.

6.10.5.1 Worked example: Real-time wildlife monitoring

The task is to design an ML system that identifies wildlife species from camera trap images in a national park. The system must process images locally with no cloud connectivity, operate on battery power for six months, and achieve 90 percent accuracy on 50 target species. The decision

process below walks through five steps: characterizing the data, analyzing the constraints, evaluating candidate architectures, validating against hardware limits, and assessing deployment risk.

The first step is data characterization. The input is spatial data (images from camera traps, typically 1920×1080 resolution, downsampled to 224×224 for processing). The task requires recognizing visual patterns (fur textures, body shapes, distinctive markings) that are:

- **Spatially local:** Species identification relies on local features (ear shape, stripe patterns)
- **Translation invariant:** A deer in the top-left is still a deer in the bottom-right
- **Hierarchical:** Low-level edges combine into textures, then body parts, then whole animals

These three properties (spatial locality, translation invariance, and hierarchical structure) point directly to a CNN, whose inductive bias matches them. Choosing a CNN captures these visual invariants natively, avoiding the parameter explosion of an unconstrained MLP.

Constraint analysis comes next. Table 6.14 catalogs the five deployment constraints and the architectural choices each forces:

Table 6.14: Wildlife Monitoring Constraints: Five deployment constraints and their implications for the camera-trap scenario. No connectivity rules out cloud-only inference, while the power, latency, memory, and accuracy requirements must be validated on the target device.

Constraint	Requirement	Implication
Connectivity	None (offline)	All inference must run on-device
Power	~2 W average (solar + battery)	Rules out GPUs; must use low-power MCU or edge NPU
Latency	<500 ms per detection	Allows batch size 1, no real-time streaming
Memory	512 MB RAM, 2 GB storage	Model, activations, and runtime buffers must fit locally
Accuracy	90%+ on 50 species	Requires sufficient model capacity

With the constraints fixed, the third step evaluates candidate architectures against the chapter's lighthouse models:

- **ResNet-50** (25.6M params, 8.2 GFLOP): Too compute- and power-heavy for this device. At 102.4 MB FP32, it is also marginal against the 100 MB this deployment allocates to model weights before lower-precision weights, activations, and runtime buffers. The main blockers are its GFLOP cost and power draw, not raw storage alone.
- **MobileNetV1** (4.2M params, 1138 MFLOP): Promising. It needs 16.8 MB at FP32, or 4.2 MB when each weight is stored as an 8-bit integer (INT8). Its depthwise separable convolutions are power-efficient.
- **KWS DS-CNN** (200K params, 20 MFLOP): Too small. Designed for 12-class audio, insufficient capacity for 50 visual species.

MobileNetV1 is therefore the right family, and its successor does the same job with fewer resources. MobileNetV2's inverted residual blocks reach comparable accuracy at a smaller parameter count, so the chosen model is a MobileNetV2 variant with width multiplier 0.75. It carries ~2.6M parameters (10.4 MB FP32, 2.6 MB with INT8 weights) and costs ~418 MFLOP at 224×224 . It is a plausible candidate for the 50-class problem; accuracy must be measured on the target data. It fits the memory budget with margin, and its depthwise separable convolutions are power-efficient.

The fourth step validates that choice against the hardware. The memory budget checks out:

$$\underbrace{2.6 \text{ MB}}_{\text{Model}} + \underbrace{224 \times 224 \times 64 \times 4}_{\text{Activations}} \approx 12.8 \text{ MB} + \underbrace{50 \text{ MB}}_{\text{OS/Buffers}} = 65.4 \text{ MB} \ll 512 \text{ MB} \checkmark$$

The compute budget holds on the target device, an ARM Cortex-A53 at 1.2 GHz with NEON SIMD (~2 GOPS INT8): $\frac{418 \times 10^6 \text{ INT8 ops}}{0.002 \times 10^{12} \text{ INT8 ops/s}} = 209 \text{ ms latency} \ll 500 \text{ ms target} \checkmark$

Power is the final check. Estimated inference power is ~200 mW for 209 ms, or 41.8 mJ per inference. At 100 inferences/day, that is 4.2 J/day before sensor, sleep, storage, and communication energy; battery-life feasibility still requires the full device budget.

The fifth step is risk assessment. Table 6.15 pairs the top accuracy, thermal, and species-coverage risks with the engineering mitigation chosen for each:

Table 6.15: Wildlife Camera Deployment Risks: Top accuracy, thermal, and species-coverage risks for the MobileNetV2 deployment, paired with the engineering mitigation chosen for each.

Risk	Mitigation
90% accuracy not achieved	Train on augmented dataset; consider EfficientNet-Lite if MobileNet insufficient
Thermal throttling in enclosure	Add passive heatsink; reduce inference frequency in high-temperature conditions
New species added postdeployment	Expand or retrain the output head and plan an over-the-air (OTA) update mechanism

The resolution is therefore MobileNetV2 (0.75× width) with INT8 weight storage, deployed on a Cortex-A53 system on chip (SoC) with 512 MB RAM. The systems insight is that this architecture fits the stated model, compute, and active-power budgets, processing images in about 209 ms under the stated throughput assumption and leaving memory headroom for system operations. We still must validate the accuracy and six-month battery targets in deployment. The decision was driven by matching the CNN inductive bias to the spatial data characteristics, then checking the hardware constraints quantitatively.

This worked example applies the chapter’s architectural principles to an engineering decision. Architecture selection still involves counterintuitive trade-offs. A model with fewer FLOPs can run slower on certain hardware. A more expressive architecture can deliver worse accuracy on problems that do not match its inductive bias. An architecture that performs well in the lab can also miss its targets on production hardware with a different memory hierarchy. The most common errors are catalogued next, each grounded in the systems principles developed throughout this chapter.

6.11 Fallacies and Pitfalls

Architecture choice is a systems decision, not a leaderboard selection. The common mistakes in this section arise when teams treat architecture families as interchangeable accuracy tools while ignoring inductive bias, memory traffic, hardware mapping, and deployment state.

Fallacy: *More complex architectures always perform better than simpler ones.*

Engineers often assume that transformers outperform simpler architectures on all tasks. In production, architectural sophistication must match problem complexity: the algorithm must fit both the structure of the data and the cost of the machine. The MNIST comparison in section 6.2.1 shows that the example CNN uses 47× fewer parameters than the example MLP because its locality bias matches image structure. Accuracy, training cost, and inference latency still require measurement for the specific implementation and task.

Pitfall: *Selecting architectures based solely on accuracy metrics without analyzing computational requirements.*

Practitioners choose architectures from papers reporting top-line accuracy, ignoring computational implications. RNNs expose a sequential critical path that limits parallelism, while transformers face quadratic dense-score computation and, in naive implementations, quadratic score storage. Sequence length 2,048 therefore requires 16× more materialized score memory than length 512. Systems that ignore these characteristics can miss latency targets, exceed memory budgets, and deliver much lower hardware utilization than expected.

Fallacy: *An architecture has one dominant bottleneck across training and inference.*

The same computation graph can occupy different systems regimes. Transformer training and prefill process many positions together and may bind on dense attention compute; small-batch autoregressive decoding repeatedly streams weights and KV-cache state and may bind on memory bandwidth. Batching, sequence length, precision, and cache layout can shift that boundary. A benchmark from one regime therefore cannot establish another’s bottleneck; measure each against the iron law.

Pitfall: *Combining architectural patterns without analyzing interaction effects at the system level.*

Engineers add attention to CNNs or convolutions to transformers expecting additive benefits. Each pattern creates distinct memory access characteristics: CNNs exploit spatial locality through sliding windows, while dense attention introduces all-to-all score computation. Combining them can

increase intermediate traffic and disrupt locality, so the hybrid’s throughput cannot be inferred by adding the components’ benefits. Adding recurrent connections to transformers likewise reintroduces sequential dependencies. Successful hybrids require profiling memory access and cache behavior before combining patterns.

Fallacy: *Architecture wins on training hardware transfer directly to deployment hardware.*

Teams design for high-end GPU clusters, then discover deployment failures on target hardware. An architecture exploiting $8 \times$ A100 GPUs (640 GB total memory) cannot deploy unchanged to a representative edge node such as the NVIDIA Jetson Orin NX (16 GB system memory). As section 6.10.3 emphasizes, architecture selection must analyze the full system stack because storage, latency, and power budgets vary by deployment target.

Pitfall: *Ignoring KV cache growth when estimating transformer serving costs.*

Teams budget transformer deployment based on model weight memory alone, overlooking the key-value (KV) cache used during autoregressive generation (Pope et al. 2023; Kwon et al. 2023). The cache scales as $\mathcal{O}(B \times N_L \times 2 \times N_{\text{heads}} \times S \times d_{\text{head}})$, where B is the number of concurrent sequences, N_L is the layer count, the factor 2 stores keys and values, N_{heads} is the head count, S is sequence length, and d_{head} is head dimension. At long contexts or high concurrency, this overhead can become a binding serving constraint. Section 6.6.4.2 works through the chapter’s 32-layer, 32-head example step by step. With 128-dimensional heads, 2,048-token sequences, and FP16 storage, each concurrent request holds ≈ 1 GB of KV cache. At 2–4 users, the cache alone consumes 2 GB–4 GB before allocator, activation, and workspace overhead. Those values do not exhaust a 80 GB device, but they consume part of the post-weight headroom and grow linearly with context and concurrency. Capacity planning must therefore budget weights, cache, and runtime buffers together, then reduce concurrency, context, or cache footprint, or add memory capacity, when the combined total approaches device memory.

6.12 Summary

Architecture is infrastructure. The choice among MLPs, CNNs, RNNs, transformers, and recommendation models (DLRM) is not merely an algorithmic hyperparameter; it fixes the physical terms in the serialized form of the iron law of ML systems ($T_{\text{exec}} = D_{\text{vol}}/\text{BW} + O/(R_{\text{peak}} \cdot \eta_{\text{hw}}) + L_{\text{lat}}$). By examining each architecture through our four-part lens (pattern processing needs, algorithmic structure, computational mapping, and system implications), we see that an architecture’s mathematical inductive bias strongly shapes which term of the iron law is likely to bind hardware performance.

The five lighthouse models established at the chapter opening map directly onto these physical bounds. ResNet-50’s spatial weight reuse inflates arithmetic intensity ($I = O/D_{\text{vol}}$), pushing execution into the compute-bound regime (O/R_{peak}). Small-batch autoregressive GPT-2 generation has low arithmetic intensity, binding execution to memory bandwidth (D_{vol}/BW). RNN temporal recurrences impose an un-parallelizable sequential critical path that raises the latency floor (L_{lat}). DLRM sparse embeddings hit a capacity wall outside the iron law: their tables can exceed accelerator memory before either execution term becomes the first feasibility test. Finally, MobileNetV2 and KWS highlight how low FLOP counts do not guarantee throughput if operational shapes fail to saturate hardware functional units.



Key Takeaways: Architecture is infrastructure

- **Inductive bias is the unifying concept:** Every architecture encodes structural assumptions: locality for CNNs, sequence for RNNs, global context for transformers. These biases trade generality for sample efficiency and determine which problems an architecture can solve efficiently.
- **Arithmetic intensity helps identify the bottleneck:** Comparing a workload’s arithmetic intensity with the target hardware’s roofline balance helps determine whether compute or memory bandwidth is likely to bind.

- **Quadratic costs are permanent constraints:** Dense transformer attention computes $\mathcal{O}(S^2)$ score interactions. Naive implementations also use quadratic score storage; tiled exact attention removes that storage cost but not the dense computation.
- **Lighthouse models isolate distinct bottlenecks:** ResNet-50 (compute), GPT-2 (bandwidth), DLRM (capacity), MobileNetV2 (latency), KWS (power). These archetypes diagnose which physical constraint dominates a given system.
- **Depth benefits from architectural support:** Skip connections and normalization improve optimization in very deep networks rather than acting as universal prerequisites beyond a fixed depth. These building blocks, born in CNNs, transfer to many deep architectures, including transformers.
- **FLOPs do not equal speed:** MobileNetV2 uses 13.7× fewer FLOPs than ResNet-50 but may run slower on some data center GPUs when its operation shapes and arithmetic intensity use the available compute units poorly. Architecture-hardware alignment, not operation count alone, determines throughput.
- **Architecture selection is deployment selection:** Choosing a transformer over a CNN determines memory requirements, latency floors, hardware utilization, and infrastructure costs. The architecture is the system constraint.


Inductive bias and systems cost are two views of the same choice. The assumptions that help a model generalize from limited samples, such as locality, recurrence, or unrestricted token interactions, also determine the shape, reuse, and lifetime of intermediate data. An algorithmic choice therefore propagates into memory capacity, bandwidth, and scheduling requirements before a framework chooses kernels.

Architecture selection signs a physical contract with hardware. A CNN commits to spatial locality and weight reuse; dense attention commits to $\mathcal{O}(S^2)$ score computation; and an RNN commits to serial time-step dependencies. These topological choices fix the fundamental workload terms O , D_{vol} , and L_{lat} , thereby governing memory footprint, bandwidth demand, and parallel efficiency. Runtime techniques can improve how efficiently the contract is paid but cannot erase it: tiled exact attention removes intermediate score storage without changing the quadratic arithmetic floor, while pipeline parallelism partitions recurrent work without removing the serial time-step path. Systems engineers can therefore anticipate likely bottleneck shifts before implementation by matching model inductive bias and data geometry to silicon physics.

What's Next: From blueprints to construction

Raw computational primitives (chapter 5) become concrete architectural contracts in this chapter. How does an abstract blueprint become an executable runtime on physical hardware? The architecture fixes *what* work must be performed; the machine learning framework determines *how* that work is scheduled and dispatched. Chapter 7 examines how PyTorch, TensorFlow, and JAX convert high-level architectural graphs into computational directed acyclic graphs (DAGs), autograd graphs, and optimized device kernels that execute directly on silicon.

ML Frameworks




7.1	Three Framework Problems
7.2	The Ladder of Abstraction
7.3	Execution Problem
7.4	Differentiation Problem
7.5	Abstraction Problem
7.6	nn.Module Abstraction
7.7	Framework Platform Analysis
7.8	Deployment Targets
7.9	Framework Selection
7.10	Anatomy of a Training Step
7.11	Fallacies and Pitfalls
7.12	Summary

Purpose

Why does the framework silently constrain every decision that comes after it?

Model equations do not execute themselves. A framework translates an architecture into work that hardware can run, and that translation is not neutral. It decides when operations execute, what intermediate state training retains, how memory is allocated, and which deployment targets can use the resulting artifact. These choices shape more than speed: they determine what can be debugged directly, what the compiler can optimize, and whether the same model survives the path from a research environment to a production runtime. The abstractions that make experimentation convenient therefore create long-lived commitments. Once a project accumulates checkpoints, modules, data loaders, serving formats, and team expertise around a framework, changing it becomes a system migration rather than a library substitution. A framework well suited to rapid exploration can expose too little structure for efficient deployment, while a constrained execution path can sacrifice the flexibility needed during discovery. Because execution, differentiation, and abstraction meet in the same software layer, a local convenience can become a downstream performance or portability constraint. The engineering question is not which framework is universally best, but which trade-offs the workload, team, and target machine can afford to make fixed. Understanding frameworks is therefore about reasoning through those system choices rather than learning transient API syntax. In D·A·M terms, the framework mediates algorithm-machine co-design by turning algorithmic intent and data representation into executable work under machine constraints.





Learning Objectives

- Explain how frameworks mediate algorithm-machine co-design through execution, differentiation, and hardware abstraction
- Compare execution strategies using dispatch overhead, compilation cost, and deployment constraints
- Analyze automatic differentiation and activation storage to choose recomputation or checkpointing
- Implement module abstraction patterns for parameter discovery, mode behavior, serialization, and hooks
- Calculate training-step FLOPs, memory traffic, and dispatch overhead to identify fusion and layout bottlenecks
- Select TensorFlow, PyTorch, JAX, or edge runtimes based on model, hardware, team, and deployment requirements

7.1 Three Framework Problems

An architecture is a graph of commitments: a transformer commits the system to attention, matrix multiplications, activation state, and memory traffic. The framework is the layer that turns that graph into work the machine can execute. A few calls such as `logits = model(tokens)`, `loss = criterion(logits, targets)`, and `loss.backward()` hide billions of floating-point operations across memory hierarchies, gradients through millions of parameters via automatic differentiation, thousands of GPU kernel launches, and gigabytes of intermediate state. The API looks simple because the framework is acting as a compiler for the silicon contract.

Architectures specify what computations neural networks perform, but knowing *what* to compute is entirely different from knowing *how* to compute it efficiently. A transformer’s attention mechanism requires coordinating computation across memory hierarchies and accelerator cores in patterns that naive implementations would execute 100× slower than optimized ones. Implementing these operations from scratch for every model would make deep learning economically infeasible. ML frameworks exist to bridge this gap by lowering abstract model graphs into hardware-aligned execution pipelines that extract maximum performance from silicon.

A framework is to machine learning what a compiler is to traditional programming. A C compiler translates human-readable code into optimized machine instructions, managing register allocation, instruction scheduling, and memory layout. An ML framework translates high-level model definitions into hardware-specific execution plans, managing operator fusion, memory reuse, and device placement. This analogy is more than metaphor: modern frameworks literally include compilers.

Every ML framework, regardless of API or design philosophy, must solve three core problems. The first is the *execution problem*: deciding when and how computation runs. A framework can execute operations immediately as written (eager execution¹) or build a complete description first—a computational graph² (a structured representation of operations and their dependencies)—and optimize before executing (graph execution). This choice shapes debugging capability, optimization potential, and deployment flexibility.

Once execution has a shape, the framework must solve the *differentiation problem*: computing gradients automatically. As established in chapter 5, training requires derivatives of a loss function with respect to millions or billions of parameters, and manual differentiation is error-prone at this scale. Frameworks therefore apply registered derivative rules across supported operation compositions while managing floating-point arithmetic and the memory overhead of saved values.

The third problem is *hardware abstraction*: targeting diverse hardware from a single interface. The same model definition should be expressible across CPUs, GPUs, Tensor Processing Units (TPUs), and mobile devices, even though each target has different memory constraints and optimal execution patterns. Within the GPU slice of this problem, some framework ecosystems expose custom-kernel languages so advanced users can write high-performance kernels without dropping all the way to low-level CUDA (Tillet et al. 2019).

¹ | **Eager execution:** This mode executes each operation immediately, which enables direct debugging with standard tools but sacrifices the global view needed for graph-level optimizations. Without capturing a larger region of computation, an eager runtime cannot fuse across that region or preplan its memory, leaving workload-dependent optimization opportunities for compilers such as `torch.compile`.

² | **Computational graph:** The “optimize before executing” distinction in the triggering sentence is the key design choice. Capturing the program as a data structure lets a framework fuse compatible operations into fewer GPU kernels before execution, reducing launch overhead and intermediate memory traffic. The engineering cost of this visibility is that the executed program differs from the source code, making debugging harder; a trade-off every graph-based framework must justify against the performance gain.

These three problems are deeply interconnected. The execution model determines *when* differentiation occurs and *what* optimizations are possible. The abstraction layer must support both execution styles across all hardware targets. Solving any one problem in isolation leads to frameworks that excel in narrow contexts but fail in broader deployment. Because these problems are ultimately about translating mathematics into efficient hardware execution, a useful perspective is to view frameworks not as libraries but as compilers.

Systems Perspective 7.1: The ML compiler

In the context of the iron law (section 1.7), a framework is a compiler for the silicon contract. The “source code” is the model architecture (the O term). The framework’s job is to take this high-level math and compile it into a series of hardware-specific kernel launches that:

1. Minimize data movement (D_{vol}) through techniques like kernel fusion.
2. Maximize utilization (η_{hw}) by matching operations to specialized hardware units like Tensor Cores.
3. Minimize overhead (L_{lat}) through efficient asynchronous dispatch and graph capture.

Choosing a framework means choosing the compiler that determines *how* efficiently a model uses hardware; precision matters: a definition that captures all three responsibilities separates genuine frameworks from numerical libraries that address only one.

Definition 7.1: Machine learning frameworks

Machine learning frameworks are software systems that translate high-level mathematical model definitions into hardware-optimized execution plans by managing the computational graph, automatic differentiation, kernel dispatch, and memory allocation across the hardware hierarchy.

1. **Significance:** Frameworks directly determine the system efficiency (η_{hw}) term in the iron law. Compiler-backed operator fusion, for example, can eliminate writes and subsequent reads of intermediate values between compatible operations. Fusing a matrix multiplication, bias add, and rectified linear unit (ReLU) changes how the same mathematics reaches hardware without changing the model.
2. **Distinction:** Unlike a numerical library such as NumPy, which normally evaluates each operation as it is called, an ML framework can capture or defer execution to analyze a larger computational graph and apply cross-operation optimizations: operator fusion, memory layout transformations, and parallel scheduling. Those optimizations require visibility beyond separately dispatched operations.
3. **Common pitfall:** A frequent misconception is that frameworks are interchangeable API wrappers. Framework choice determines which compiler paths are available. PyTorch can recover graph-level optimization from eager code through `torch.compile()`, while TensorFlow and JAX commonly rely on XLA-backed compilation paths that lower operations to target hardware. Moving from eager execution to a compiler path can change throughput materially, but the gain depends on model structure, shapes, operator support, and hardware.

The compiler metaphor is not decorative. An ML framework translates logical intent into physical execution under the constraints of the iron law, deciding how to partition computation across memory hierarchies, when to trade numerical precision for throughput, and how to schedule operations so that the dominant term (data movement, computation, or overhead) is minimized. The framework is where the governing physics developed throughout this book becomes executable code.

The scale of this translation is not obvious from the API surface. A single call to `loss.backward()` triggers operation recording, memory allocation for gradients, reverse-order graph traversal, and

hardware-optimized kernel dispatch—machinery that would require hundreds of lines of manual calculus for even a three-layer network. For a contemporary language model, the framework additionally orchestrates billions of floating-point operations across accelerators, coordinating memory hierarchies, numerical precision, and, when the system grows beyond one device, communication libraries. Building this from scratch would be economically prohibitive for most organizations, which is why the history of ML frameworks is a history of progressively automating these layers.

The three problems—execution, differentiation, and abstraction—did not emerge simultaneously. Each arose as a response to scaling limitations in the previous generation of tools. Tracing this evolution explains why modern frameworks are designed as they are and why they embody these particular trade-offs.

7.2 The Ladder of Abstraction

In 1979, writing a matrix multiplication in Fortran that used the hardware efficiently required deep knowledge of cache lines, register scheduling, and vector units. By 2016, a single line of Python (`torch.matmul(A, W)`) could dispatch to a highly optimized implementation without the programmer knowing the details of the silicon. That compression of effort did not happen in one step; it accumulated across four decades of abstraction, each layer solving a bottleneck that made the previous generation impractical for scaling. The result is a **ladder of abstraction** where each rung automates what the rung below exposed.

1. **Solving Performance (1979–1992):** The original **Basic linear algebra subprograms (BLAS)**³ standardized reusable low-level linear-algebra primitives (Lawson et al. 1979), while LAPACK⁴ (Bai et al. 2006) built higher-level numerical routines on top of that foundation. Together, these libraries solved the problem of hardware primitives: stable interfaces let frameworks delegate operations such as $C = A @ B$ ⁵ to specialized implementations instead of hand-writing silicon-specific code.
2. **Solving Usability (2005–2006):** NumPy⁶ solved the problem of developer velocity. By wrapping low-level BLAS routines in high-level Python (Harris et al. 2020), it allowed scientists to write code in a friendly language while executing it in optimized C/Fortran. This “Vectorization” pattern, where the slow language handles logic and the fast language handles loops, became a durable contract for scientific computing. Jupyter notebooks later extended this usability layer into readable, executable computational workflows for combining code, results, and explanations (Kluyver et al. 2016).
3. **Solving Differentiation (2007–present):** Deep learning frameworks (Theano,⁷ TensorFlow (Abadi et al. 2016), PyTorch (Paszke et al. 2019)) solved the problem of gradient computation. While NumPy required manual derivation of backpropagation gradients (error-prone and slow), these frameworks made automatic differentiation through computational graphs a standard capability. This turned the chain rule into a software primitive, allowing researchers to define *forward* passes and get *backward* passes automatically.

Machine learning frameworks evolved by progressively abstracting hardware execution details into higher-level API primitives. Frameworks bridge the gap between mathematical intent and silicon reality (figure 7.1), moving numerical software up the ladder from low-level linear algebra routines to compiled graph execution engines.

Each generation abstracted away details that consumed engineering effort in the previous one, yet each abstraction introduced new trade-offs. BLAS hid assembly-level optimization but fixed the interface. NumPy hid memory management but required manual differentiation. Modern frameworks hide gradient computation but introduce a choice among execution models. The deeper pattern is that every abstraction hides details by preserving a contract about shape, dtype, device, and sometimes units; when that contract becomes implicit, correct components can still compose into a wrong system.

With that contract risk in view, modern frameworks converge on the same three core problems: *how* to execute computation, *how* to differentiate it, and *how* to abstract across hardware. The execution problem comes first because its resolution determines what optimizations the other two problems can exploit.

³ | **BLAS (basic linear algebra subprograms):** The 1979 API specification that forms the bottom rung of the ladder described here; it standardized a fixed set of Fortran-callable vector operations, separating the public routine names from the machine-specific implementation decisions beneath them. Every framework above it inherits the broader version of this bargain: call a standard linear-algebra primitive from any language and let a tuned vendor library target the silicon. For modern general matrix multiply (GEMM) on NVIDIA GPUs, that tuned path is cuBLAS rather than the 1979 BLAS specification itself (NVIDIA 2024a).

⁴ | **LAPACK (linear algebra package):** Extends BLAS by providing a standard API for higher-level routines (SVD, eigendecomposition, least-squares) that vendors implement with chip-specific code layered on top of fast GEMM kernels. This layered design is the architectural pattern every ML framework inherits: high-level operations delegate downward to hand-tuned primitives, so a vendor-optimized LAPACK call can execute over 10× faster than a naive implementation without the framework author writing a single line of hardware-specific code.

⁵ | **GEMM:** The matrix-matrix primitive behind $C = A @ B$. Hardware vendors hand-tune GEMM for their specific chips because dense layers, attention projections, and convolution lowering all rely on matrix multiplication, making this one routine a performance floor for many frameworks above it on the ladder. Its high arithmetic intensity makes GEMM the operation most able to approach peak compute throughput, while small or misaligned shapes often fall back to much lower utilization.




Figure 7.1: Computational Library Evolution: A timeline of numerical computing milestones spanning 1979 to 2018. Each era automates what the layer below exposed—from BLAS and LAPACK hardware routines to NumPy vectorized syntax and graph-based automatic differentiation in modern frameworks—trading low-level transparency for developer velocity.

⚠ War Story 7.1: The interface that forgot its units (1999)

Context: NASA’s Mars Climate Orbiter depended on ground software, spacecraft software, navigation teams, and contractor interfaces all agreeing on the physical meaning of exchanged values (Mars Climate Orbiter Mishap Investigation Board 1999).

Mechanism: Ground software supplied by Lockheed Martin computed thruster impulse in English units (pound-force seconds), while JPL navigation software consumed the values assuming SI units (newton-seconds)—a $4.45\times$ implicit unit mismatch.

Impact: On September 23, 1999, the spacecraft was lost after entering occultation. Its corrected 57 km periapsis, rather than the planned 226 km, was judged too low for survival.

Response: The board recommended unit-consistency checks, specification audits of transferred data, and stronger end-to-end verification of ground software.

Systems lesson: Framework abstractions are valuable because they carry contracts: shape, dtype, device placement, and physical units. If those contracts are implicit, two pieces of correct code can still compose into a wrong system. ML frameworks hit this whenever tensor memory layouts (such as NCHW vs. NHWC), quantization scale factors, or physical feature units differ across pipeline stages, producing plausible outputs that quietly corrupt downstream inference.

7.3 Execution Problem

Consider two engineers writing the same neural network. The first debugs interactively, printing tensor shapes after each operation, inspecting intermediate values, and stepping through code with `pdb`. The second waits 30 seconds for compilation, then watches the model run $3\times$ faster while losing that direct, line-by-line view of intermediate state. Both are correct; they have made different choices about the **execution problem**, the question of whether operations should execute immediately as written or be recorded for later execution. This choice creates a cascade of engineering trade-offs that shape every aspect of framework behavior, from debugging workflows to deployment options to peak hardware utilization.

7.3.1 Why execution strategy matters: The memory wall

To understand why execution strategy matters so much, return to the widening compute-bandwidth gap quantified by the memory-wall equation (equation 2.3). Processor arithmetic has grown faster than memory bandwidth, creating the memory wall. Modern accelerators can perform arithmetic far faster than they can fetch data. Element-wise operations like ReLU use only a tiny fraction of peak compute capacity, not because the hardware is slow, but because they spend nearly all their time waiting for data. Section D.2.1 formalizes this trade-off, showing exactly when operations are memory bound vs. compute bound.

The memory wall classifies operations as either compute-bound (limited by arithmetic throughput) or memory-bound (limited by data movement). Most individual neural network operation types (activations, normalizations, element-wise operations) are memory bound, though the large matrix multiplications that dominate total compute time can be compute bound.

6 | NumPy (numerical Python): In 2005, Travis Oliphant unified two competing Python array libraries (Numeric and Numarray) into a single package, giving the scientific computing community one BLAS-backed array standard at the moment it needed to scale. The “vectorization” contract this created (write logic in Python, execute loops in C/Fortran via BLAS) became the design template for every ML framework that followed: PyTorch tensors and TensorFlow arrays are direct descendants, extending the same n -dimensional array abstraction to GPUs. Python’s role in ML infrastructure inherits much of its shape from this consolidation decision.

7 | Theano: Developed at the Montreal Institute for Learning Algorithms (MILA) under Yoshua Bengio starting in 2007, Theano was an early and influential Python framework that compiled symbolic mathematical expressions into optimized CPU and GPU code via computational graphs (Bergstra et al. 2010; Team et al. 2016). It demonstrated that a Python-defined computation graph could be compiled for GPU execution without requiring the researcher to write CUDA code. Mila ended active development of Theano in 2017 (Mila 2026).

8 | Kernel (GPU): In GPU programming, a kernel is the function dispatched to execute in parallel across thousands of threads. Each kernel launch incurs 5–20 μ s of CPU-side overhead for parameter assembly and GPU signaling, which means that small, unfused operations spend more time on launch overhead (L_{lat}) than on useful arithmetic. Reducing kernel count through fusion is therefore a direct attack on the overhead term of the iron law.

9 | Fused attention kernel: A fused attention kernel combines the QK^T product, softmax, and value-weighted output into a tiled implementation that keeps intermediate values in on-chip memory rather than materializing the full attention matrix in HBM. FlashAttention is the canonical named example introduced in chapter 6 and reports up to $9\times$ fewer HBM accesses with workload-specific speedups from 15 percent to $3\times$ (T. Dao et al. 2022). The framework lesson is broader than the specific algorithm: fusion can shift an operation’s position on the Roofline Model from bandwidth-limited toward throughput-limited execution by reducing round-trips through external memory.

The key optimization for memory-bound operations is kernel fusion, combining multiple operations into a single GPU function (called a *kernel*)⁸ to avoid intermediate memory traffic. Fusing a sequence of normalization, dropout, and activation operations into one kernel can yield large speedups by eliminating intermediate writes between operations. Attention kernels⁹ use the same principle at larger scale: instead of materializing the full attention matrix in high-bandwidth memory (HBM), a fused implementation can keep tiles close to the compute units, cut HBM accesses by up to $9\times$, and produce workload-specific speedups from 15 percent to $3\times$ (T. Dao et al. 2022).

Frameworks can fuse only operations visible together. Separately dispatched eager operations hide cross-operation opportunities, but graph capture can expose them. A deferred graph lets the framework optimize the captured computation. Execution strategy therefore determines which optimizations are possible and their scope.

7.3.2 The computational graph

Kernel fusion is the key optimization for memory-bound operations, but fusion requires seeing multiple operations together. Frameworks make this visibility possible through the computational graph, a directed acyclic graph (DAG) where nodes represent operations and edges represent data dependencies. This graph is the framework’s internal model of the computation.

Mathematical operations can be decoupled from physical execution through graph representations. The computational graph (figure 7.2) grounds this abstraction: tensor variables map to data nodes while operations map to transformation nodes.




Figure 7.2: Simple Computational Graph: Representing $z = f(x, y)$ as a directed acyclic graph decouples mathematical dependencies from physical execution order, providing the structural blueprint frameworks require to track gradients and schedule operations across devices.

Real machine learning models require much more complex graph structures. Figure 7.3 extends this representation to show a neural network computation graph alongside the system components that reason about it. In the left panel, notice how data flows through six operation nodes in a directed acyclic graph—each node’s output becomes the next node’s input. The right panel reveals what the framework gains from this explicit graph structure: it can query the structure to plan memory allocation for each tensor’s lifetime, and it can assign operations to devices based on data dependencies rather than execution order. The critical insight is that the graph exists independently of execution, enabling the framework to optimize before any arithmetic occurs.




Figure 7.3: Computation Graph with System Interactions: A neural network DAG (left) interfaces with framework runtime components (right). By maintaining an explicit graph representation independent of execution, the framework can analyze tensor lifetimes for memory reuse and assign operations to target accelerators before issuing arithmetic.

This graph representation is more than a visualization; it is the data structure that enables both efficient execution and automatic differentiation. The answer to *when* this graph is constructed creates a design choice with cascading implications across four dimensions. Debugging benefits from visibility into intermediate values and step-through execution. Optimization benefits from seeing multiple operations at once, which enables fusion. Deployment benefits when execution no longer depends on the Python interpreter. Flexibility benefits when control flow can depend on computed tensor values.

No single execution model optimizes all these dimensions. Frameworks must choose their position in this trade-off space, and practitioners must understand these trade-offs to select appropriate tools and write efficient code. The three execution families that follow are different answers to the same systems question: how much graph visibility should the framework trade for immediate execution and debugging?

7.3.3 Three execution strategies

The computational graph representation enables global optimization, but it leaves a critical design choice unresolved: *when* the framework builds the graph. Consider a simple operation like $y = x * 2$. One approach performs the multiplication immediately, storing the result in y . This is natural and debuggable, but the framework sees only one operation at a time. The other approach defers execution, recording the intention to multiply and building a graph of operations that runs later when explicitly requested. This is less intuitive, but the framework sees the complete computation, which enables optimization.

Neither approach dominates; each embodies different trade-offs between *flexibility* and *optimization potential*. Modern frameworks have explored three primary execution strategies: eager execution with dynamic graphs, static computation graphs, and hybrid approaches that combine just-in-time (JIT) compilation with eager development. Each strategy has distinct systems implications.

7.3.3.1 Eager execution with dynamic graphs

Eager execution evaluates each operation immediately as the program calls it, building the computation graph dynamically at runtime. A side-by-side comparison shows how this differs from graph-based execution at the code level.

Example 7.1: Eager vs. graph execution code comparison

PyTorch (eager execution):

```
import torch

x = torch.tensor([1.0, 2.0])
y = x * 2
print(f"Intermediate value: {y}") # Works immediately
z = y.sum()
```

TensorFlow 1.x (static graph):

```
import tensorflow as tf

x = tf.placeholder(tf.float32)
y = x * 2
# print(y) -> Prints Tensor("mul:0"...), not value!
z = tf.reduce_sum(y)

with tf.Session() as sess:
    result = sess.run(z, feed_dict={x: [1.0, 2.0]})
```

Systems insight: Eager execution exposes intermediate values as ordinary runtime state, which makes debugging direct. Static graphs stage computation before execution, which enables whole-graph optimization but changes the debugging model.

Eager execution runs operations immediately as encountered, building the computation graph dynamically during execution. When a programmer writes $y = x * 2$, the multiplication happens instantly and the result is available for immediate use.

This provides the flexibility of normal programming: developers can print intermediate values, use conditionals based on computed results, and debug with standard tools. The framework records operations as they happen, constructing a dynamic graph that reflects the actual execution path taken.

For gradient computation, the framework records a history of operations in what is called an autograd tape,¹⁰ a transient data structure built during execution. Each tensor operation creates a node that records: the operation performed, references to input tensors, and how to compute gradients. These nodes form a DAG that records the actual path taken during forward-pass execution rather than a graph fixed in advance. Listing 7.1 shows how PyTorch records operations as they execute in its default eager mode.

Listing 7.1: Autograd Tape Construction: Each operation executes immediately while recording a backward node to the autograd tape for later gradient computation.

```
import torch

x = torch.tensor([1.0], requires_grad=True)
y = x * 2 # Executes immediately; records MulBackward node
z = y + 1 # Executes immediately; records AddBackward node
# The autograd tape exists NOW, built during execution
```

After these two operations, the framework has constructed an autograd tape with two nodes: one for the multiplication and one for the addition. The tape records that z depends on y , and y depends on x .

Calling `z.backward()` traverses this tape in reverse topological order, applying the chain rule at each node:

1. Compute $\frac{\partial z}{\partial z} = 1$ (seed gradient)
2. Call `AddBackward0.backward()` $\rightarrow \frac{\partial z}{\partial y} = 1$
3. Call `MulBackward0.backward()` $\rightarrow \frac{\partial z}{\partial x} = 2$
4. Accumulate gradient in `x.grad`

After `backward()` completes, the autograd graph is normally released. The next forward pass builds a new graph. Values needed for gradients are saved during the forward pass and remain live until backward consumes them, so their memory cost spans both phases rather than appearing only during backward.



Example 7.2: In-place operations can break gradients

Scenario: An engineer applies an in-place mutation (`x += 1`) within a custom PyTorch activation function to reduce memory allocations.

Diagnosis: In-place operations overwrite tensor memory containing forward-pass activations required by the autograd tape for backward-pass gradient computation, triggering a runtime version-counter error (PyTorch Contributors 2026a).

Systems lesson: Framework automatic differentiation depends on immutable forward-pass activation records. Unchecked in-place memory mutations break autograd tape dependencies, forcing runtime framework execution halts.

In eager execution frameworks, operations evaluate imperatively as host code encounters them. Trace the execution loop in figure 7.4, following the define-and-dispatch sequence that sends individual operations directly to device stream queues.

¹⁰ | **Autograd tape:** A transient data structure built during forward execution, where nodes record operations, dependencies, and backward functions for chain-rule evaluation. Its memory footprint grows with model depth and sequence length as saved activations accumulate until released during backpropagation. For deep models, frameworks retain selected activations and recompute others via activation checkpointing to prevent OOM failures.




Figure 7.4: Dynamic Graph Execution Flow: In eager execution, each operation is issued as the program encounters it. Accelerator work may execute asynchronously after dispatch. This define-by-run model enables natural debugging and data-dependent control flow while limiting optimization across Python calls.

Systems implications: Flexibility The dynamic autograd tape expresses data-dependent behavior directly in the host language. Conditionals and loops can depend on tensor values computed during execution, enabling algorithms like beam search, dynamic recurrent neural network lengths, or adaptive computation that adjust their behavior based on intermediate results. Static graphs can also represent data-dependent control flow and dynamic dimensions through graph operations, but eager execution makes these patterns natural to write and debug in Python. Because operations are issued immediately, developers can print tensors, inspect values, and use standard debuggers (pdb, breakpoints) to diagnose errors much as they would in another Python program.

Systems implications: Overhead This flexibility comes with performance costs that map directly to the iron law (section 1.7). Each training forward pass constructs an autograd tape dynamically, adding host-side Python dispatch overhead and graph-recording work to L_{lat} . Operations pass through the framework dispatcher before device kernels are launched, so overhead becomes significant when individual kernels are short. Eager execution alone also lacks a static globally captured graph, preventing the framework compiler from performing kernel fusion (combining multiple element-wise operations into a single GPU kernel launch to eliminate intermediate HBM reads/writes) or pre-allocating a static memory arena for intermediate activations (D_{vol}). The autograd tape retains the values requested by backward rules, adding memory pressure until those values can be released. Together, these costs create a performance ceiling that becomes visible as operations grow smaller and dispatch overhead dominates computation.

The dispatch tax: Python overhead vs. GPU reality Eager execution’s performance ceiling is driven by a fundamental systems mismatch: the speed of the host-side interpreter vs. the speed of the device-side silicon. The dispatch tax, defined as the fraction of time spent in the host-side orchestration (Python) vs. actual device execution (GPU), quantifies this mismatch.


In the illustrative eager-execution model used here, every operation pays a fixed “tax” of approximately $15\ \mu s$ for host-side dispatch, type checks, and kernel launch. Actual overhead varies by framework, backend, and hardware, but its relative weight still depends on operation size. For a small operation such as a ReLU on a small vector, the kernel might execute in only $1\ \mu s$, so the modeled dispatch tax reaches 94 percent and the GPU spends most of its time waiting for the next command. For a large operation such as a large matrix multiply, the kernel executes for $100\ \mu s$, the modeled dispatch tax drops to 13 percent, and the system becomes compute bound.

The dispatch tax explains why models with many small layers run significantly slower than their raw FLOP count predicts. Section A.5.1 places this symptom in the bottleneck taxonomy, classifying a dispatch-dominated workload as latency-bound rather than compute-bound and showing which optimizations actually move it. To approach efficient execution, frameworks must move from *kernel-by-kernel dispatch* to *graph-level execution*, where the dispatch tax is paid once for the entire graph rather than per operation. The hybrid JIT and compilation strategies in section 7.3.3.3 exist precisely to address this overhead.

The overhead costs of eager execution motivate the opposite design: capturing the entire computation before executing any of it. This is precisely what static computation graphs provide.

7.3.3.2 Static computation graphs

Static graph execution defines the complete computational graph as a symbolic representation first, then executes it separately. This “define-then-run” execution model means the graph exists before any computation occurs, enabling aggressive ahead-of-time optimization. The key insight is that if the framework sees the entire computation before running it, the framework can analyze, transform,



As work grows, useful compute outpaces fixed dispatch cost; the tax shrinks.

and optimize the graph globally—visibility unavailable across separate eager calls unless a compiler captures them. Section C.3.4.2 works through the canonical global transformation this enables: for N elements of b bytes each, fusing a chain of k elementwise operations into one kernel reduces idealized read-and-write traffic from $2kNb$ to $2Nb$ bytes.

Two-phase execution Static graphs implement a clear separation between graph construction and execution. Listing 7.2 illustrates the two phases using TensorFlow 1.x, which exemplified this approach. It deliberately runs the same $x * 2$ then $+ 1$ computation shown under eager execution in listing 7.1, holding the arithmetic fixed so that the only thing that changes is *when* it executes: symbolic definition creates placeholders and operations without computation, while explicit execution triggers actual arithmetic.

Listing 7.2: Static Graph Two-Phase Execution: Graph construction (symbolic definition) is separated from execution (actual computation), enabling ahead-of-time optimization.

```
# Phase 1: Graph Construction (symbolic, no computation)
import tensorflow.compat.v1 as tf

tf.disable_v2_behavior()

# Define graph symbolically
x = tf.placeholder(tf.float32, shape=[1]) # Just a placeholder
y = x * 2 # Not executed, just recorded
z = y + 1 # Still no execution
# At this point, nothing has been computed

# Phase 2: Graph Execution (actual computation)
with tf.Session() as sess:
    result = sess.run(z, feed_dict={x: [1.0]})
    # Now computation happens: result = [3.0]
```

Static graph compilation separates graph construction from tensor execution. Observe the phase division in figure 7.5, noting how symbolic graph definition on the left precedes ahead-of-time optimization and runtime execution on the right.




Figure 7.5: Static Graph Execution Lifecycle: Separating symbolic graph definition (left) from concrete execution (right) provides global visibility over the computation, allowing compilers to apply ahead-of-time optimizations—such as operator fusion, buffer reuse, and layout transforms—before any tensors are materialized.

The key difference from eager execution is that during construction, x , y , and z are not tensors containing values but rather symbolic nodes in a graph. Operations like $*$ and $+$ add nodes to the graph definition without performing any arithmetic. The `print(y)` line in the code example would reveal this distinction—it would print tensor metadata, not a computed value. Execution is triggered explicitly through `sess.run()`, at which point the framework analyzes the complete graph, optimizes it, and executes the optimized version with the provided input data.

Ahead-of-time optimization Because the framework has the complete graph before execution, it can perform ahead-of-time optimization [optimizing the graph before runtime] unavailable across uncaptured eager calls. The kernel fusion opportunity introduced in section 7.3.1 becomes actionable here: because the framework sees $y = x * 2$ and $z = y + 1$ together in the graph, it can fuse them into $z = x * 2 + 1$, eliminating the intermediate y and halving idealized memory traffic. When tensor shapes and lifetimes are known, the compiler can plan memory before execution and reuse buffers whose lifetimes do not overlap; dynamic dimensions may require bounds or runtime allocation. Tensor layouts can also be transformed globally (for example, NCHW to NHWC) to

match hardware preferences, though physical layout conversions may still require copies. Dead-code elimination (DCE)¹¹ removes operations whose results are unused, and constant folding precomputes operations on constant values when their inputs are known. These optimizations map directly to iron law terms: kernel fusion can reduce D_{vol} by avoiding intermediate memory writes, constant folding reduces repeated work, planned allocation can reduce runtime overhead, and dead code elimination removes work and associated data movement when unused computations are present.

Compilation frameworks like **XLA (accelerated linear algebra)**¹² (Google 2025) take this further, compiling TensorFlow graphs to optimized executables for specific hardware.

Systems implications Static graphs can achieve high performance through ahead-of-time optimization. Kernel fusion reduces memory bandwidth requirements when intermediate traffic is the bottleneck, and hardware-specific compilation can approach high utilization when shapes, operators, and layouts match efficient backend paths.

The cost of this performance is reduced flexibility. Standard Python control flow (`if`, `for`) cannot depend on computed tensor values in static graphs. TensorFlow provides graph-level control flow primitives (`tf.cond` and `tf.while_loop`) that support data-dependent conditions, but these require special syntax that diverges from standard Python, making code harder to write and reason about. Debugging is difficult because stack traces point to graph construction code, not execution code. Error messages often reference symbolic node names rather than the actual operations that failed.

7.3.3.3 Hybrid approaches: JIT compilation

JIT compilation pursues both eager debugging and graph optimization at once by capturing computation at runtime. The core trade-off is *fidelity vs. generality*. Tracing captures the execution path taken during a sample run, producing high fidelity to that path but missing branches not taken. Source-level compilation (scripting) analyzes supported program structure, preserving supported control flow but requiring a restricted language subset. Both approaches produce an intermediate representation (IR)¹³ that enables graph optimizations such as operator fusion, constant folding, dead code elimination, and buffer reuse.

The eager-vs.-compiled trade-off has a direct iron law consequence. JIT compilation amortizes the L_{lat} (dispatch overhead) across the compiled region. Longer compiled regions mean more overhead amortized per operation, which explains why graph breaks are performance-critical: each break forces a return to eager dispatch, resetting the amortization.

Historically, PyTorch’s TorchScript exemplified both strategies (PyTorch Contributors 2026d). TorchScript is now deprecated in favor of newer capture and export paths such as `torch.export` (PyTorch Contributors 2026d, 2026c); it remains useful here as a concrete illustration of tracing and source-level scripting. Tracing executes a function with example inputs and records the tensor operations observed on that path. Listing 7.3 demonstrates how a traced module could be serialized and executed independently of the Python interpreter.

Listing 7.3: TorchScript Tracing: Captures tensor operations by executing a function with example inputs and recording the execution path into a static computation graph.

```
import torch

def forward(x):
    y = x * 2
    z = y + 1
    return z

# Trace the function by running it once
x_example = torch.tensor([1.0])
traced = torch.jit.trace(forward, x_example)

# traced is now a compiled TorchScript module
# Can serialize: torch.jit.save(traced, "model.pt")
# Can optimize: fusion, constant folding
# Can run without Python interpreter
```

¹¹ | **Dead code elimination (DCE):** Removes graph nodes whose results do not affect observable outputs or side effects. In ML graphs, dead code arises from unused branches or post-transform artifacts; side-effect analysis ensures safe removal to prevent execution of dead operations and reduce kernel launch overhead.

¹² | **XLA (accelerated linear algebra):** The “optimized machine code” in the triggering sentence means XLA can fuse subgraphs, specialize layouts, and lower high-level operations into backend-specific code; fusion attacks both launch overhead (L_{lat}) and intermediate memory writes (D_{vol}), but the realized speedup depends on whether the graph contains enough fusible, memory-bound work for the compiler to remove. Large GEMM-heavy regions may already be compute bound, while chains of small elementwise operations can benefit more because fusion removes repeated trips through external memory. The benefit is not a fixed multiplier: XLA helps when it can fuse operations, specialize layouts and shapes, and reduce launch or memory overhead, so gains depend on graph structure, backend support, and input-shape stability.

¹³ | **Intermediate representation (IR):** The “intermediate” captures this format’s architectural role: a language-independent layer that decouples the frontend (Python capture) from the backend (hardware code generation), exactly as LLVM IR decouples C/Rust/Swift frontends from x86/ARM backends. ML frameworks adopted this compiler pattern because it reduces the $\mathcal{O}(M \times N)$ cost of supporting M frontends and N backends to $\mathcal{O}(M + N)$: a single graph capture mechanism (TorchDynamo, tf2xla) can target multiple hardware backends without rewriting the capture logic.

The critical limitation of tracing reveals the fidelity-generality trade-off concretely. A trace records operations from an observed path rather than preserving arbitrary Python control flow. Listing 7.4 illustrates the resulting correctness risk.

Listing 7.4: Tracing a Data-Dependent Branch: Tracing records the path taken by the example input. PyTorch commonly emits a `TracerWarning` when a tensor value is converted to a Python condition because the resulting trace may not generalize.

```
def conditional_forward(x):
    if x.sum() > 0: # Data-dependent condition
        return x * 2
    else:
        return x * 3

traced = torch.jit.trace(conditional_forward, torch.tensor([1.0]))
# Tracing captures ONLY the x.sum() > 0 branch
# If input later has sum <= 0, traced version
# still executes x * 2 branch
```

Tracing records the branch executed by the example input. In this example, converting the tensor condition to a Python Boolean commonly produces a `TracerWarning`; the resulting trace can still follow the captured branch for later inputs that require the other branch. Treating trace warnings as correctness failures is therefore essential when data-dependent Python control flow is present.

TorchScript’s historical alternative, scripting, analyzed supported Python source directly and compiled it to TorchScript IR without tracing a sample path (PyTorch Contributors 2026d). The scripting compiler preserved supported branching structure but accepted only a restricted Python subset. Current projects should follow the supported `torch.export` and target-runtime guidance rather than select between TorchScript tracing and scripting for new deployment pipelines (PyTorch Contributors 2026c). Within the historical TorchScript workflow, tracing fit feed-forward models whose tensor paths remained stable for the supported inputs, while scripting handled supported data-dependent control flow without a Python interpreter. Scripting’s key advantage was its ability to preserve supported conditionals in the IR, as listing 7.5 shows.

Listing 7.5: Scripted Control Flow: Unlike tracing, scripting preserves both branches of conditionals in the IR, enabling correct execution based on runtime input values.

```
@torch.jit.script
def conditional_forward(x: torch.Tensor) -> torch.Tensor:
    if x.sum() > 0:
        return x * 2
    else:
        return x * 3

# Both branches preserved in IR
# Correct branch executes based on runtime input values
```

To understand what the compiler produces, listing 7.6 inspects the generated intermediate representation directly, where the single Python expression has been lowered into explicit typed primitive operations the runtime can execute without the interpreter.

Scripting imposed constraints because TorchScript had to analyze code that Python normally interprets dynamically. Function signatures and variables sometimes needed type annotations, and unsupported Python objects, library calls, or metaprogramming could fail compilation. Table 7.1 summarizes this historical design trade-off rather than current deployment guidance.

The TorchScript IR represents operations using the `aten` namespace for core tensor operations, the `prim` namespace for primitives and control flow, static types for every value, and static single-assignment (SSA) form, where each variable is assigned exactly once to simplify compiler analysis. This IR enables optimizations independent of Python: operator fusion combines adjacent operations

into single kernels, constant folding evaluates constant expressions at compile time, dead code elimination removes unused operations, and memory optimization reuses buffers when possible.

Listing 7.6: TorchScript IR Inspection: The generated intermediate representation shows primitive operations and constants, useful for debugging and understanding compilation results.

```
@torch.jit.script
def example(x: torch.Tensor) -> torch.Tensor:
    return x * 2 + 1

# Inspect generated IR:
print(example.graph)
# graph(%x : Tensor):
#   %1 : int = prim::Constant[value=2]()
#   %2 : Tensor = aten::mul(%x, %1)
#   %3 : int = prim::Constant[value=1]()
#   %4 : Tensor = aten::add(%2, %3, %3)
#   return (%4)
```

Table 7.1: Historical TorchScript Tracing vs. Scripting Trade-Offs: Tracing recorded operations from example executions, while scripting preserved supported control flow at the cost of a restricted Python subset. TorchScript is deprecated for new projects.

Aspect	Tracing	Scripting
Input requirement	Example inputs needed	No inputs needed
Control flow	Cannot handle data-dependent	Supports data-dependent
Conversion ease	Simpler (just run function)	Harder (restricted Python)
Type annotations	Not required	Required when inference fails
Error detection	Runtime (wrong results)	Compile time (syntax errors)
Best for	Feed-forward models	Models with conditionals

7.3.3.4 Modern compilation: Graph-capture JIT

The previous approaches force a choice between flexible eager execution and compiler-visible graphs. Modern JIT compilation narrows this trade-off by automatically capturing regions of eager code into optimized graphs with limited developer intervention.

Graph-capture JIT systems follow the same architectural pattern across frameworks. The first execution observes tensor operations, records a graph region guarded by assumptions about shapes, dtypes, layouts, and control flow, lowers that region into an intermediate representation, applies fusion and layout optimizations, and caches executable code for later calls that satisfy the same guards. Unsupported Python code does not disappear; it forms a graph boundary where execution returns to the eager runtime. PyTorch 2.0’s `torch.compile` (Ansel et al. 2024) is a concrete instance of this pattern, but the systems idea is broader than the API: compilation pays only when captured regions are long enough and stable enough to amortize capture, lowering, code generation, and cache-management costs.

This explains when compilation can help. Dispatch costs that seem negligible for one operation—a few microseconds at a time—accumulate across the thousands of operations in a forward pass. A simple fusion estimate makes the overhead concrete.

 Napkin Math 7.1: The physics of software overhead

Iron law connection: The latency term (L_{lat}) in the iron law is dominated by software overhead: dispatching instructions from Python to the GPU. Each operation pays a Python dispatch cost of $\sim 10\text{ }\mu\text{s}$ plus a kernel launch cost of $\sim 5\text{ }\mu\text{s}$, which together set the per-launch overhead applied in the scenarios that follow.

Scenario one: Eager Mode (The “Tiny Op” Trap) Consider a simple activation block, $y = \text{relu}(x + \text{bias})$, in which each tensor contains N elements of b bytes each.

- **Operations:** Two (Add, ReLU).
- **Execution:**
 1. Launch Add Kernel: $15 \mu\text{s}$ overhead.
 2. Read/Write Memory: $2Nb$ bytes.
 3. Launch ReLU Kernel: $15 \mu\text{s}$ overhead.
 4. Read/Write Memory: $2Nb$ bytes.
- **Total overhead:** $30 \mu\text{s}$.
- **Total memory traffic:** $4Nb$ bytes.

Scenario two: Compiled Mode (Fusion) The compiler fuses this into one kernel: FusedAddRelu.

- **Execution:**
 1. Launch Fused Kernel: $15 \mu\text{s}$ overhead.
 2. Read/Write Memory: $2Nb$ bytes (the intermediate result stays on chip).
- **Total overhead:** $15 \mu\text{s}$ ($2\times$ speedup).
- **Total memory traffic:** $2Nb$ bytes ($2\times$ bandwidth efficiency).

Systems insight: Fusion wins on two fronts at once. Collapsing two launches into one halves the per-op dispatch overhead in this scenario, and keeping the intermediate result on chip cuts idealized memory traffic from $4Nb$ to $2Nb$ bytes. For small element-wise operations, the avoided round-trip to external memory can matter more than the arithmetic.

¹⁴ | **torch.compile:** It is a 2020s PyTorch implementation of graph-capture JIT compilation: bytecode interception extracts tensor regions from eager programs, an intermediate representation carries those regions to compiler backends, and cached generated code is reused while guard conditions continue to hold.

¹⁵ | **CUDA (compute unified device architecture):** NVIDIA’s parallel computing platform (2007) serving as the foundational layer between high-level Python operations and GPU silicon; when PyTorch executes `torch.matmul(A, W)`, the call traverses the framework’s dispatcher, selects a cuBLAS kernel, and launches it on the GPU. Each launch incurs $5\text{--}20 \mu\text{s}$ of CPU-side overhead. For small operations, this dispatch overhead (L_{lat}) exceeds the useful compute time, which is why compilation (fusing N_{ops} operations into one kernel launch) yields speedups proportional to the reduction in launch count rather than the reduction in arithmetic.

Figure 7.6 makes the dispatch tax visible: eager execution creates gaps where the GPU sits idle while Python dispatches the next kernel. The blue compute regions are short; the red dispatch regions are comparatively long. Compilation fuses these operations into a single kernel launch, replacing many dispatch gaps with one dispatch block and one fused compute block.




Figure 7.6: The Python Tax: Execution timelines show how eager mode incurs repeated host dispatch latency (L_{lat}) between short GPU kernels, leaving the accelerator periodically idle. Compilation fuses these operations into a single kernel launch, eliminating intermediate dispatch gaps and maximizing hardware occupancy.

Automating this fusion is the design goal behind graph-capture compilers such as PyTorch 2.0’s `torch.compile`.¹⁴ They capture eager tensor regions and compile them into fused kernels without requiring engineers to write custom CUDA.¹⁵

The core engineering questions are therefore conceptual, not API-specific. A capture compiler needs a frontend that identifies graph regions in an eager program, an intermediate representation that separates the captured computation from Python, and a backend that lowers the region to hardware-specific code. It also needs a guard system: the compiled artifact is valid only while assumptions about tensor rank, dtype, layout, and control-flow path remain true. When a guard fails, the runtime must recompile or fall back to eager execution.

Graph breaks mark the boundary where compilation stops applying. Data-dependent Python control flow, unsupported library calls, I/O, custom Python objects, and highly variable shapes all shorten compiled regions. Each break reintroduces dispatch overhead and may require tensors to move between compiled code and the eager runtime. This is why graph-break analysis belongs in performance engineering: the relevant metric is not whether compilation is enabled, but how much of the hot path remains inside long, stable compiled regions.

Backends occupy different points on the flexibility-performance spectrum. A **JIT** compiler traces model execution at runtime to generate optimized kernels dynamically, incurring initial warmup latency during the first pass. An **ahead-of-time (AOT)** compiler instead compiles captured computation before deployment, shifting compilation out of the hot path and enabling a standalone runtime when the export target supports one. A general JIT backend optimizes ordinary training and serving workloads with moderate compilation cost; a specialized inference backend can apply deeper fusion, precision lowering, and autotuning when the deployment target is fixed; an ahead-of-time mobile or embedded runtime removes even more flexibility to gain footprint and predictability. The same rule governs all of them: the narrower the target and the more stable the graph, the more optimization the compiler can safely perform.

The resulting workflow is a systems decision. Rapid prototyping favors eager execution because architecture changes and guard failures make recompilation cost visible. Long training runs and high-volume inference amortize compilation cost over many executions, provided the model has stable shapes and limited graph breaks. Debugging usually starts in eager mode because errors map directly to source code; compilation is reintroduced after the model behavior is correct and the performance bottleneck is measurable.

Comparison of execution models Table 7.2 contrasts the three execution models across six dimensions, showing how hybrid JIT compilation can recover graph-level optimization within captured regions while preserving eager fallback outside them.

Table 7.2: Execution Model Trade-Offs: Each execution model occupies a distinct position in the flexibility-optimization trade-off space. Eager execution maximizes debugging flexibility but limits cross-operation optimization; static graphs maximize optimization but require control flow to be represented in the graph; hybrid JIT compilation compiles captured regions while falling back to eager execution for unsupported patterns.

Aspect	Eager + Autograd Tape (PyTorch default)	Static graph (TensorFlow 1.x)	JIT Compilation (torch.compile)
Execution Model	Immediate	Deferred	Hybrid
Graph Construction	During forward pass	Before execution	First execution (cached)
Optimization	Per-operation kernels	Ahead-of-time	JIT compilation
Dynamic Control Flow	Python control flow	Graph control flow	Captured regions or breaks
Debugging	Easy (standard Python)	Difficult (symbolic)	Moderate (mixed)
Performance	Baseline	High (optimized)	High (compiled regions)

Eager mode’s primary value is in the iteration loop quantified in section 3.1: it allows using standard Python debuggers (like `pdb`) to inspect variables mid-execution, whereas graph-mode debugging often requires specialized framework tools. This immediate feedback accelerates the prototyping phase of the ML lifecycle.

Beyond these core execution trade-offs, table 7.3 highlights additional systems-level distinctions between static and dynamic approaches.

These trade-offs are not binary choices. Modern frameworks offer a spectrum of options, which raises the quantitative question of where on this spectrum a given project should operate.

Table 7.3: Additional Graph Trade-Offs: Systems-level distinctions between static and dynamic graphs that complement table 7.2. These trade-offs reappear when selecting frameworks in section 7.9.

Aspect	Static Graphs	Dynamic Graphs
Memory Management	Precise allocation planning, optimized memory usage	Flexible but potentially less efficient
Hardware Utilization	Can generate highly optimized hardware-specific code	May sacrifice hardware-specific optimizations
Research Velocity	Slower iteration due to define-then-run requirement	Faster prototyping and model experimentation
Integration with Legacy Code	More separation between definition and execution	Natural integration with imperative code

7.3.4 Quantitative principles of execution

These execution models form a spectrum, and two quantitative principles make the choice measurable in practice. The compilation continuum principle asks when compilation gains justify development cost by comparing production executions with development iterations. The dispatch overhead law shows why eager execution can spend more time dispatching small operations than computing them.

7.3.4.1 The compilation continuum principle

The execution problem demands a quantitative principle for when a project should compile. The execution models form a continuum from maximum flexibility to maximum optimization. Equation 7.1 lays out the four positions on that axis, and each labeled arrow names the mechanism that carries a project one step rightward toward hardware.

$$\text{Eager} \xrightarrow{\text{capture}} \text{Graph JIT} \xrightarrow{\text{export}} \text{AOT Runtime} \xrightarrow{\text{specialize}} \text{Hardware} \quad (7.1)$$

Each step rightward sacrifices flexibility for performance. The practical question is *where* a given project should be placed on this continuum. The optimal compilation strategy depends on production executions, development recompilations, and their respective time costs, combined in equation 7.2:

$$\text{Compilation Benefit} = \frac{N_{\text{prod}} \cdot (T_{\text{eager}} - T_{\text{compiled}})}{T_{\text{compile}} + N_{\text{dev}} \cdot T_{\text{compile}}} \quad (7.2)$$

Where:

- N_{prod} = number of production executions (dimensionless count: inference requests, training steps)
- N_{dev} = number of development iterations requiring recompilation (dimensionless count)
- T_{eager} = time per execution in eager mode (seconds)
- T_{compiled} = time per execution in compiled mode (seconds)
- T_{compile} = time per compilation or recompilation (seconds; assumed constant)

This model assumes that the compiled artifact remains valid between development changes and that the initial compilation and each recompilation have the same cost; under these assumptions, the decision rule is to compile when $\text{Compilation Benefit} > 1$. The ratio is dimensionless.

Table 7.4 presents a hypothetical throughput scenario across execution modes and model architectures:

Table 7.4: Illustrative Training and Inference Throughput: Hypothetical throughput and compilation-time inputs for comparing execution modes on an NVIDIA A100 GPU with batch size 32. These values are assumptions rather than benchmark measurements. Within this scenario, the table implies `torch.compile` speedups of 1.4–1.5× over eager mode, while TensorRT implies 2.3–2.7× speedups but requires longer compilation and is inference only.

Model	Eager (examples/sec)	<code>torch.compile</code> (examples/sec)	TensorRT (examples/sec)	Compile Time (seconds)
ResNet-50	1,450	2,150	3,800	15–30
BERT-Base	380	520	890	30–60

Model	Eager (examples/sec)	torch.compile (examples/sec)	TensorRT (examples/sec)	Compile Time (seconds)
ViT-B/16	620	950	1,650	25–45
GPT-2 (124M)	180	260	420	45–90

These throughput differences across execution modes raise a practical question—which framework execution strategy best serves each workload archetype. The appropriate strategy depends on the workload, backend, and dominant iron law term, and table 7.5 maps each recurring archetype to a candidate execution strategy.

Table 7.5: Framework Execution Strategy by Workload: Candidate execution strategies for each workload archetype, aligned to the dominant iron law term.

Archetype	Dominant Iron Law Term	Candidate Framework Strategy	Rationale
ResNet-50 (Compute Beast)	$\frac{O}{R_{\text{peak}} \eta_{\text{hw}}}$ (Compute)	Compiled dense kernels	Regular dense kernels benefit from layout selection, precision lowering, and backend specialization; fusion helps most in surrounding memory- or launch-bound regions
GPT-2 (Bandwidth Hog)	$\frac{D_{\text{vol}}}{\text{BW}}$ (Memory Bandwidth)	Fused attention + graph compilation	Fused attention and compilation reduce HBM round-trips and improve cache reuse
DLRM (Sparse Scatter)	$\frac{D_{\text{vol}}}{\text{BW}_{\text{random}}} + L_{\text{lat, network}}$	Eager execution with specialized kernels	Embedding lookups are inherently irregular and dynamic; compilation gains are small
DS-CNN (Tiny Constraint)	L_{lat} (Overhead)	Ahead-of-time microcontroller runtime	Sub-ms inference; every microsecond of Python overhead is unacceptable



Lighthouse 7.1: Framework strategy by archetype

The four recurring archetypes in table 7.5 trace one principle across the framework: compilation benefits scale with how much of the workload is optimizable.

Systems insight: In the illustrative scenario, Compute Beasts such as the ResNet-50 row in table 7.4 benefit because their dense kernels expose substantial surface for layout selection, precision lowering, and fusion. Sparse Scatter workloads such as DLRM can gain less when irregular embedding lookups leave limited compiler scope.

This principle has concrete implications across three regimes. In research prototyping ($N_{\text{dev}} \gg N_{\text{prod}}$), teams should stay eager. If the architecture changes every few minutes, compilation overhead can dominate. Under the scenario’s 30-second compilation assumption, ten compilations per hour consume five minutes.

For long training runs ($N_{\text{prod}} \gg N_{\text{dev}}$), compilation can pay off because its one-time cost is amortized across many steps. In the hypothetical ResNet-50 scenario in table 7.4, torch.compile provides 48.3 percent higher throughput (2,150 img/s vs. 1,450 img/s). Using the scenario’s 30 s compile cost, this pays off after the breakeven point in equation 7.3:

$$N_{\text{breakeven}} = \frac{T_{\text{compile}}}{T_{\text{eager}} - T_{\text{compiled}}} \quad (7.3)$$

Evaluating equation 7.3 with the hypothetical ResNet-50 values gives approximately 134,000 images, well within a single training run of any realistic length.

For production inference ($N_{\text{dev}} \approx 0$, $N_{\text{prod}} \rightarrow \infty$), teams should maximize compilation. With no development iterations and potentially millions of requests, every optimization matters. Aggressive autotuning can be worthwhile even when compilation takes much longer, because the cost is amortized over the deployment lifetime.

These three regimes create distinct regions in the compilation decision space. Figure 7.7 maps out these regions so engineers can identify where each strategy wins. Watch for the crossover points: the steep eager line (highest per-execution cost) eventually overtakes JIT's moderate slope, while the gentlest compiled line (lowest per-execution cost but largest upfront investment) wins only beyond the second crossover in this illustrative model. The slopes reveal per-execution cost; the vertical offsets reveal compilation overhead. A project's execution volume and measured costs determine which strategy minimizes total time.




Figure 7.7: The Compilation Continuum: An illustrative cost model shows how one-time compilation cost and per-execution cost create eager, JIT, and AOT crossover regions as production executions increase. The crossover locations depend on the assumed compilation costs and per-execution rates.

7.3.4.2 The dispatch overhead law

A second principle, the dispatch overhead law, emerges from equation 7.4, which identifies the regime in which framework overhead, rather than compute or memory, dominates execution time. Let N_{ops} be the number of operations (count), $t_{dispatch}$ the per-operation dispatch overhead (seconds), and $T_{compute}$ and T_{memory} the total compute and memory times (seconds). The ratio uses the additive, no-overlap approximation $T_{hw} \approx T_{compute} + T_{memory}$; when execution overlaps those costs, measured hardware time should replace their sum. Framework overhead dominates when operations are small relative to dispatch cost:

$$\text{Overhead Ratio} = \frac{N_{ops} \cdot t_{dispatch}}{T_{compute} + T_{memory}} \quad (7.4)$$

When Overhead Ratio > 1 , the model is overhead-bound. Compilation can provide its largest gains for overhead-bound workloads because fusion reduces the number of dispatches. The training-step trace in section 7.10 works this effect through end to end, previewed by the numbers that follow.

Summed over N_{ops} operations, the per-operation dispatch cost $t_{dispatch}$ accumulates into a per-call tax on execution. Whether that tax dominates depends on how the hardware execution time T_{hw} compares to the software overhead T_{sw} (both measured in seconds), and the regimes split sharply by model size.

The principle's implication is that workloads composed of small operations can benefit disproportionately from compilation when dispatch dominates execution. Model parameter count alone does not determine the gain.

The dispatch tax analysis reveals that small operations become overhead-bound when dispatch time exceeds device execution time. This observation matters most at the extreme edge of the deployment spectrum, where the Python runtime itself may exceed the target's resource budget.

Napkin Math 7.2: The dispatch tax

Problem: When does Python overhead kill performance?

Scenario one: Small multilayer perceptron (MLP) (Overhead Bound)

- **Compute:** 6 ops across small matrix/element-wise operations.
- **Hardware time:** $T_{hw} \approx 2.6 \mu s$ (mostly memory latency).
- **Software overhead:** $T_{sw} \approx 6 \text{ ops} \times 15 \mu s/\text{op} = 90 \mu s$.
- **Ratio:** $90 \mu s / 2.6 \mu s \approx 34.6$.
- **Small-model outcome:** The system spends 97 percent of time in host-side dispatch and kernel-launch overhead. Eliminating all modeled dispatch overhead would give an upper-bound speedup of 35.6×

Scenario two: GPT-3 Layer (Compute Bound)

- **Compute:** Huge matrix multiplications.
- **Hardware time:** $T_{hw} \approx 100 \text{ ms} = 100000 \mu s$.
- **Software overhead:** $T_{sw} \approx 50 \mu s$.
- **Ratio:** $50 \mu s / 100000 \mu s \approx 0.0005$.
- **Large-model outcome:** Modeled Python overhead is negligible. Compilation may still improve fusion, layout, or backend specialization; dispatch elimination contributes little.

Systems insight: Dispatch overhead is regime-dependent. Compilation can reduce host-side overhead for small-operation workloads by fusing operations into fewer launches, while large models benefit mainly from fused kernels and memory movement reductions.

7.3.5 Frameworks for the edge: TinyML and micro-runtimes

The compilation continuum reaches its extreme at the far edge. While cloud frameworks like PyTorch and TensorFlow 2.x prioritize flexibility through eager execution, TinyML¹⁶ systems operating on MCUs with kilobytes of memory cannot afford the overhead of a Python interpreter or a fully dynamic runtime.

Lighthouse 7.2: Lighthouse example: KWS on TinyML

Scenario: Deploying the Smart Doorbell’s keyword spotting (KWS) model, the DS-CNN Tiny Constraint lighthouse, to an ARM Cortex-M4 with 256 KB RAM and 1 MB Flash.

Constraint: A full PyTorch and Python runtime exceeds the device’s memory budget by orders of magnitude.

Framework solution: Micro-frameworks such as **TensorFlow Lite Micro (TFLM)** (David et al. 2021) solve this through a tiny interpreter-based runtime with a fixed memory discipline:

1. **Fixed memory arena:** The application supplies a contiguous tensor arena, and the framework plans and reuses buffers from that arena rather than relying on dynamic allocation during inference.
2. **Kernel selection:** Only the specific kernels used by the model (for example, Conv2D, DepthwiseConv) need to be linked or registered with the runtime.
3. **Compact interpreter execution:** The MCU runs a small C/C++ interpreter over a flat model representation, with the model and arena bound at initialization rather than assembled dynamically at runtime.

Silicon contract: On TinyML devices, the contract is strictly memory-bound. The framework’s primary job is to ensure the model’s intermediate activations (the “working set”) fit within the MCU’s tiny SRAM.

¹⁶ | **TinyML:** Systems designed for MCUs cannot afford a Python interpreter. Micro-runtimes instead use small C/C++ runtimes, fixed memory planning, and model-specific operators. TensorFlow Lite Micro remains interpreter-based: the application supplies a tensor arena whose buffers are planned and reused without heap allocation after setup. Predictable memory is mandatory because exceeding a 256 KB device budget can prevent initialization or inference.

These micro-runtimes sacrifice dynamic flexibility to reduce footprint and make memory predictable. With compact models, low-power hardware, and duty cycling, that discipline enables milliwatt-scale inference with mostly on-chip data movement.

The spectrum of execution strategies, from dynamic eager execution to static graph compilation and specialized micro-runtimes, requires developers to make deliberate trade-offs. The key execution-mode decisions summarize these architectural choices:

Checkpoint 7.1: Execution models

The choice of execution mode determines both developer velocity and model performance.

Debuggability vs. speed

- ☐ **Eager mode (Python-first):** why does executing ops one-by-one make debugging easy but optimization hard?
- ☐ **Graph mode (compiler-first):** why does building a static graph enable kernel fusion?

Modern compromise

- ☐ **JIT compilation:** how does `torch.compile` bridge the gap?
- ☐ **Dispatch-overhead estimate:** if six operations each incur the chapter's 15 μs /op dispatch tax while hardware work takes 2.6 μs , which term dominates?

The execution problem determines *when* computation happens and *what* optimizations are possible. Neural network training, however, requires a capability that no amount of clever scheduling can provide: the ability to compute gradients automatically.

Consider what training actually requires: for each of millions of parameters, compute how a tiny change would affect the loss. Doing this manually for even a simple three-layer network requires deriving and implementing dozens of partial derivatives. For a modern transformer with billions of parameters, manual differentiation is economically impossible. A framework that executes efficiently but cannot differentiate can run inference but cannot learn.

7.4 Differentiation Problem

The **differentiation problem** is the task of computing gradients¹⁷ automatically. Neural network training requires derivatives of a scalar loss $\mathcal{L}$ with respect to millions or billions of parameters, making manual differentiation impractical. Because a single scalar loss depends on many parameters, reverse-mode automatic differentiation (AD)¹⁸ is normally efficient: one reverse traversal computes all requested parameter gradients, while recovering the same full gradient with forward mode requires one tangent direction per parameter. Major ML frameworks therefore use reverse-mode AD for ordinary scalar-loss training (Baydin et al. 2018).

Building on the backpropagation algorithm introduced in chapter 5, the systems engineering of differentiation addresses how frameworks represent computation graphs, manage memory for intermediate values, and orchestrate the backward pass efficiently across accelerators. The framework's role is not to perform calculus but to manage the bookkeeping at scale, which is required for the training algorithms detailed in chapter 8. Listing 7.7 illustrates the core idea with a simple three-operation function.

Listing 7.7: Automatic Differentiation: AD decomposes complex functions into elementary operations with known local derivatives, allowing frameworks to apply the chain rule over a recorded operation graph.

```
def f(x):
    a = x * x # Square
    b = sin(x) # Sine
    c = a * b # Product
    return c
```

Frameworks decompose this function into elementary operations, each with a known local derivative, and then combine these local derivatives via the chain rule to compute gradients through

¹⁷ | **Automatic differentiation (AD):** The “automatically” in the triggering sentence is the key word: AD mechanizes the chain rule as a graph traversal, eliminating the manual derivative computation that made scaling beyond toy networks impractical. The systems trade-off that makes this feasible is the choice of reverse mode, which exploits the many-to-one topology of training (many parameters, one scalar loss) to compute all gradients in a single backward pass. Forward mode would require one pass per parameter, making billion-parameter training computationally impossible.

¹⁸ | **Reverse-mode AD:** The $\mathcal{O}(1)$ -vs.- $\mathcal{O}(P)$ asymmetry of reverse mode has a concrete price: reverse mode retains the operation-dependent values required by derivative rules during the backward traversal. Activation checkpointing changes this saved set by omitting selected values and recomputing them during the backward pass.

supported compositions. The systems challenge is implementing this efficiently: the framework must record the computation graph during the forward pass, store intermediate values, and execute the backward pass with minimal memory overhead. Production frameworks solve each of these problems through structured graph recording, intermediate state management, and reverse-mode traversal.

7.4.1 Forward and reverse mode differentiation

Two primary approaches to automatic differentiation exist, and the choice between them (forward mode vs. reverse mode) determines whether gradient computation scales with the number of input or output directions. This distinction explains why neural network training usually relies on reverse mode. Forward mode is useful to understand first because it exposes the bookkeeping directly; reverse mode then appears as the systems response to the many-parameter, one-loss shape of neural network training.

7.4.1.1 Forward mode

Neural network training usually uses reverse mode (covered next), but forward mode illuminates why reverse mode is efficient for a scalar loss with many parameters. Forward mode automatic differentiation computes directional derivatives alongside the original computation, tracking how changes propagate from input to output. This approach mirrors manual derivative computation, making it intuitive to understand and implement.

Forward mode propagates a tangent alongside each primal value and does not need to retain a reverse tape for a later backward traversal. Its memory and operation overhead depend on the computation and on how many tangent directions are propagated. One seeded execution produces a Jacobian-vector product (JVP). Recovering a full gradient with respect to P independent inputs requires P tangent directions, which may be evaluated separately or batched but still scales with the input dimension. Reverse mode instead computes a scalar loss gradient in one backward traversal whose cost depends on the operations rather than the parameter count alone. This asymmetry makes reverse mode the standard choice for neural-network training, while forward-mode JVPs remain useful when the number of input directions is small and in higher-order differentiation methods.

To see the mechanism concretely, consider computing both the value and derivative of $f(x) = x^2 \sin(x)$. Listing 7.8 shows how forward mode propagates derivative computations alongside every operation, applying the chain rule and product rule at each step:

Listing 7.8: Forward Mode AD: Propagates derivatives forward through the computation graph, computing one directional derivative per seeded forward pass.

```
def f(x): # Computing both value and derivative
    # Step 1:  $x \rightarrow x^2$ 
    a = x * x # Value:  $x^2$ 
    da = 2 * x # Derivative:  $2x$ 

    # Step 2:  $x \rightarrow \sin(x)$ 
    b = sin(x) # Value:  $\sin(x)$ 
    db = cos(x) # Derivative:  $\cos(x)$ 

    # Step 3: Combine using product rule
    result = a * b # Value:  $x^2 * \sin(x)$ 
    dresult = a * db + b * da # Derivative:  $x^2 * \cos(x) + \sin(x) * 2x$ 

    return result, dresult
```

Forward mode achieves this systematic derivative computation by augmenting each number with its derivative value, creating what mathematicians call a “dual number”. Listing 7.9 runs the same function at $x = 2.0$, so the bookkeeping becomes concrete: each intermediate value carries its derivative alongside it through a single pass.

Listing 7.9: Dual-Number Trace: The forward-mode function of listing 7.8 evaluated at $x = 2.0$. Each line keeps a value and its derivative, so one pass yields both $f(2.0) \approx 3.637$ and $f'(2.0) \approx 1.973$.

```
x, dx = 2.0, 1.0 # seed: track the derivative with respect to x

a = x * x # value: 4.0
da = 2 * x # derivative: 4.0
b = sin(x) # value: 0.9093
db = cos(x) # derivative: -0.4161
c = a * b # value: 3.637
dc = a * db + b * da # derivative: 1.973
```

That paired execution propagates one seeded tangent direction alongside the primal values. Its overhead depends on the operations and implementation rather than being exactly $2\times$. For a scalar loss with P parameters, recovering the full gradient through forward mode requires P independent tangent directions, which may be evaluated separately or batched but still scales with the input dimension.

Forward mode is therefore inefficient for the usual training shape of one scalar loss and millions of parameters. It remains useful when the number of seeded input directions is small and in mixed-mode methods for specialized derivatives such as Jacobian-vector and Hessian-vector products. For the usual scalar-loss training objective, reverse mode reverses this scaling: one backward traversal computes the requested gradients for all parameters that influence the loss.

7.4.1.2 Reverse mode

Modern ML frameworks normally use reverse mode for scalar-loss training because of computational asymmetry. Forward mode propagates derivatives for selected input directions, while one reverse traversal propagates adjoints from the scalar loss to all requested parameters. The advantage grows with the number of parameters, although the actual cost also depends on the graph and derivative rules.

This asymmetry makes reverse mode the standard choice for neural network training, while forward and mixed modes remain useful for other derivative shapes. Reverse mode runs on the same three-operation function $f(x) = x^2 \sin(x)$ from listing 7.7, where x reaches the output through two distinct paths: the square and the sine. Algorithm 7.1 states the general reverse-mode contract before a worked trace makes one concrete input visible.

Algorithm 7.1 Reverse-mode automatic differentiation

Require: directed acyclic graph with scalar output y ; nodes $v_1, \dots, v_n$ in topological order; a local derivative rule for each node

Ensure: adjoints $\bar{v}_i = \partial y / \partial v_i$ for the requested input or parameter nodes

- 1: **for** $i = 1$ to n **do**
 - 2: compute v_i , record dependencies, and save values requested by its derivative rule $\triangleright$ *forward pass*
 - 3: **end for**
 - 4: $\bar{v}_i \leftarrow 0$ for all i ; $\bar{y} \leftarrow 1$
 - 5: **for** $i = n$ down to 1 **do**
 - 6: **for** each parent u of v_i **do**
 - 7: $\bar{u} \leftarrow \bar{u} + \bar{v}_i \partial v_i / \partial u \triangleright$ *accumulate via the chain rule*
 - 8: **end for**
 - 9: release stored values once every rule that needs them has run
 - 10: **end for**
 - 11: **return** the adjoints for the requested nodes
-

The structure of algorithm 7.1 is also its cost. The forward pass records graph dependencies and retains values requested by backward rules until the reverse traversal consumes them. Reverse mode therefore trades activation and graph memory for an efficient gradient of a scalar output. Checkpointing can save fewer values and recompute them later.

For the concrete function in listing 7.7 with $x = 2.0$, we record $a = x^2 = 4.0$, $b = \sin(x) \approx 0.9093$, and $c = ab \approx 3.637$ during the forward pass. Seeding $\bar{c} = 1.0$, the reverse multiplication gives $\bar{a} = 0.9093$ and $\bar{b} = 4.0$; the square path contributes $0.9093 \cdot 4.0 \approx 3.6372$ to $\bar{x}$, and the sine path adds $4.0 \cos(2.0) \approx -1.6646$. The final derivative is $\partial c / \partial x = \bar{x} \approx 1.973$.

The critical observation is that this single backward pass computed $\partial c / \partial x$ regardless of how many paths connect x to c . In a neural network, each weight can affect the loss through thousands of paths across layers, and reverse mode handles them all in one traversal. This is why training a 175B-parameter model like GPT-3 is feasible at all: reverse mode's $\mathcal{O}(1)$ backward passes relative to parameter count keep gradient computation tractable.

Translating this mathematical elegance into a working system requires solving a concrete engineering problem: the backward pass needs values computed during the forward pass, so the framework must decide what to store, when to store it, and when to free it. Modern frameworks accomplish this through computational graphs and automatic gradient accumulation.¹⁹

Listing 7.10 illustrates this with a two-layer network, showing both the forward computation that stores intermediate values and the backward pass that consumes them to produce gradients for every parameter simultaneously.

Listing 7.10: Reverse Mode in a Neural Network: The forward pass computes and stores intermediate values; the backward pass walks the computation in reverse to produce gradients for every parameter.

```
def simple_network(x, w1, w2):
    hidden = x * w1 # First layer
    activated = max(0, hidden) # ReLU activation
    output = activated * w2 # Second layer
    return output

# --- Forward pass stores intermediates ---
x, w1, w2 = 1.0, 2.0, 3.0
hidden = x * w1
activated = max(0, hidden)
output = activated * w2

# --- Backward pass consumes them ---
d_output = 1.0 # Seed gradient
d_w2 = activated # = 2.0
d_activated = w2 # = 3.0
d_hidden = d_activated * (1 if hidden > 0 else 0) # ReLU gate: 3.0
d_w1 = x * d_hidden # = 3.0
d_x = w1 * d_hidden # = 6.0
```

Three implementation requirements emerge from this example. First, the framework must track dependencies between operations to determine the correct reverse traversal order. Second, intermediate values (hidden, activated) must persist in memory until the backward pass consumes them. Third, every operation needs both a forward implementation and a corresponding backward rule. These requirements define the engineering surface of any AD system, and the second requirement, memory persistence, turns out to be the dominant cost.

7.4.1.3 Memory management strategies

A 175B-parameter model in FP16 requires 350 GB just for weights, beyond a single GPU's memory. Reverse mode AD also saves operation-dependent forward values for backward. In an illustrative 100-layer network with a batch of 64 images, activations alone can add 8–12 GB alongside weights, gradients, and optimizer state. Memory capacity can therefore set the feasibility boundary even when compute is available.

The saved-activation footprint accumulates across layers. Listing 7.11 shows how each added layer can contribute another activation tensor that must persist until the backward pass reaches it.

¹⁹ | Gradient accumulation:

The framework processes a large logical batch as smaller mini-batches, freeing each mini-batch's saved activations after backward while accumulating parameter-sized gradients. Processing 4,096 samples in 64-sample mini-batches can reduce the batch-dependent activation peak by up to 64× when activation memory scales linearly with batch size; weights, gradients, optimizer state, and other fixed costs do not shrink. The trade-off is more sequential work before each optimizer update.

Listing 7.11: Activation Persistence Across Layers: The marked intermediates represent values retained for backward; peak saved-activation memory depends on their number and sizes.

```
def deep_network(x, w1, w2, w3):
    # Forward pass - must store intermediates
    hidden1 = x * w1
    activated1 = max(0, hidden1) # Store for backward
    hidden2 = activated1 * w2
    activated2 = max(0, hidden2) # Store for backward
    output = activated2 * w3
    return output
```

Frameworks attack this memory wall with two primary strategies. The first is activation checkpointing (also called gradient checkpointing): rather than storing every activation, the framework keeps selected boundary values and recomputes the missing intermediates during the backward pass. At this point, the important systems idea is the runtime contract, not the placement policy. The framework treats some activations as durable checkpoints and treats the rest as values that can be regenerated when the backward traversal reaches them; section 8.6.5.1 later examines how training systems choose the checkpoints. Listing 7.12 makes the runtime contract visible: the forward pass keeps only the segment boundaries, and the backward pass re-runs each segment to regenerate the activations it dropped rather than reading them back from memory.

Listing 7.12: Activation Checkpointing: The forward pass stores only segment boundaries; the backward pass re-runs each segment to regenerate the dropped activations, trading recomputation for memory.

```
# Standard backward: every forward activation stays resident
h1 = layer1(x) # kept for backward
h2 = layer2(h1) # kept for backward
out = layer3(h2) # kept for backward

# Checkpointed: keep only the boundary h1 and drop h2. The
# backward pass re-runs the wrapped segment to regenerate
# h2 on demand instead of holding it in memory.
h1 = layer1(x) # boundary: kept
out = checkpoint(
    lambda a: layer3(layer2(a)), h1
) # h2 recomputed in backward
```

²⁰ | **Operation fusion:** When compatible operations execute as separate kernels, intermediate results may be written to HBM and read back by the next kernel. Fusion can keep those values on chip and avoid the corresponding transfers. The realized benefit depends on the operations, shapes, compiler, and hardware.

The second strategy is *operation fusion*.²⁰ Rather than executing matrix multiplication, bias addition, and ReLU as three separate operations that materialize intermediate results, frameworks can fuse compatible work into a single kernel. This can accelerate memory-bound portions by keeping intermediate values in registers or caches.

The backward pass itself benefits from hardware-specific optimization. Rather than directly translating the mathematical definition of a convolution gradient into code, frameworks implement specialized backward kernels that exploit memory access patterns and hardware capabilities of modern accelerators (Chetlur et al. 2014). These optimizations, checkpointing, fusion, and specialized kernels, work together to make training practical for architectures that would otherwise exhaust GPU memory in a single forward pass.

7.4.2 Framework implementation of automatic differentiation

Checkpointing, fusion, and specialized kernels address the systems problems of AD. Frameworks usually expose these mechanisms through high-level APIs. A PyTorch training loop—`optimizer.zero_grad()`, forward pass, `loss.backward()`, `optimizer.step()`—appears to be four function calls. Behind each call, however, the framework tracks operations during the forward pass, builds and maintains the computation graph, manages memory for intermediate values, schedules gradient computations, and interfaces with hardware accelerators. The same graph machinery extends to advanced scenarios: nested `torch.autograd.grad` calls compute second-order derivatives for

techniques like natural gradient descent, and mixed-precision contexts (`autocast`) select reduced-precision kernels for compute-intensive operations while maintaining FP32 for numerical stability.

7.4.2.1 PyTorch autograd internals

The autograd system is the framework component that solves the differentiation problem described in section 7.1. Three systems principles govern its design: the data structure that enables efficient gradient computation, the memory cost of maintaining that data structure, and the control mechanisms that production systems require. Understanding these principles explains why training can consume substantially more memory than weight-only inference for the same model, and why frameworks provide specific mechanisms to manage that cost.

The reverse-linked graph structure During the forward pass, the autograd system constructs a reverse-linked graph of Function nodes. Each node records the operation performed and stores references to the tensors it needs for gradient computation. This graph is the data structure that makes reverse-mode automatic differentiation possible: regardless of how many parameters a model has, a single backward pass through this graph computes all gradients. For a model with P parameters, reverse-mode AD requires $\mathcal{O}(1)$ backward passes (compared to $\mathcal{O}(P)$ for forward-mode), which is why every major framework implements this approach.

Concretely, every tensor produced by a differentiable operation stores a `grad_fn` attribute pointing to the Function that created it. Each Function links to its inputs through `next_functions`, forming a chain from the loss back to the leaf parameters. Listing 7.13 illustrates this structure for a simple computation:

Listing 7.13: Reverse-Linked Graph Structure: Each tensor's `grad_fn` links to the Function that created it, forming a reverse chain from output to leaf parameters that enables $\mathcal{O}(1)$ backward passes.

```
import torch

x = torch.tensor([2.0], requires_grad=True)
y = x * 3
z = y.pow(2)


# Traverse the reverse-linked graph
print(z.grad_fn) # PowBackward0
print(z.grad_fn.next_functions) # -> MulBackward0
print(
    z.grad_fn.next_functions[0][0].next_functions
) # -> AccumulateGrad (leaf)
```

The traversal reveals the chain: `PowBackward0` (for `z = y**2`) links to `MulBackward0` (for `y = x * 3`), which terminates at `AccumulateGrad` for the leaf tensor `x`. Leaf tensors are the endpoints of the graph where gradients accumulate into the `.grad` attribute rather than propagating further. The tuple format (Function, index) tracks which output of a multi-output operation each connection corresponds to.

A reverse-linked autograd structure has a critical systems implication: the entire graph must remain in memory from the time a tensor is created until the backward pass consumes it. The graph itself is lightweight (pointers and metadata), but the tensors it references are not, so memory consumption scales with model depth.

The memory-compute trade-off Every value saved for the backward pass persists until its derivative rule consumes it. These saved values are one major reason training uses more memory than inference, alongside gradients and optimizer state. Different operations save different inputs or outputs, and checkpointing can replace some storage with recomputation.

Consider a sample ResNet-50 training scenario with 25.6M parameters (~102.4 MB of FP32 weights), batch size 64, and 224×224 images. With 8 GB–12 GB of saved activations, gradients add ~102.4 MB, and two Adaptive Moment Estimation (Adam) FP32 moment buffers add ~204.8 MB. Summing those four components gives an 8.4 GB–12.4 GB training scenario, roughly $82.1\text{--}121.2\times$ the FP32



weight storage alone. It is not a comparison with complete inference memory, which also includes activations and runtime workspaces.

The example shows why training capacity cannot be estimated from weights alone. Both depend on the operations, implementation, and memory reuse, as derived in section C.3.3 for the four-component training-state equation ($M_{\text{total}} = M_{\text{weights}} + M_{\text{gradients}} + M_{\text{optimizer}} + M_{\text{activations}}$) and additional runtime memory.

Frameworks provide three primary mechanisms to manage this trade-off at the graph level. Gradient checkpointing (T. Chen et al. 2016) changes what the graph preserves: instead of saving all activations, the framework saves selected boundary values and rebuilds the missing intermediates during the backward pass. In iron law terms, checkpointing increases the O term (recomputation) to reduce the D_{vol} term (memory traffic). Tensor detachment provides a complementary mechanism: calling `.detach()` on a tensor changes which graph edges participate in differentiation, preventing the framework from saving activations through that path. This is essential for transfer learning, where pretrained layers should not accumulate gradients, and it reduces the D_{vol} term by eliminating unnecessary activation storage. Mixed-precision training offers a third approach: store selected activations and matrix operations in lower-precision formats so the framework reduces data movement while preserving numerically sensitive work in FP32. Chapter 8 develops these graph mechanisms into sizing decisions.

Extensibility and control Production training systems require fine-grained control over gradient flow that goes beyond the default backward pass. Three categories of control arise in practice. First, **selective gradient computation**: transfer learning and fine-tuning require freezing subsets of parameters, which the framework supports through `requires_grad=False` flags and the `.detach()` mechanism described earlier. Second, **gradient inspection and modification**: debugging vanishing or exploding gradients, implementing per-tensor gradient clipping, and logging gradient statistics all require intercepting gradients mid-computation, which frameworks expose through hook APIs. Third, **custom differentiation rules**: operations not in the framework’s built-in library (custom CUDA kernels, novel activation functions, domain-specific operations) require user-defined forward and backward implementations.

These control mechanisms share a common systems design: they are callback-based extensions that the autograd engine invokes at specific points during graph traversal, without modifying the core differentiation algorithm. This extensibility pattern allows the framework to maintain a single optimized backward pass while supporting arbitrarily complex gradient manipulation. In practice, this control appears through a few recurring PyTorch mechanisms: retained graphs, accumulated gradients, custom backward rules, hooks, and safe detachment. Table 7.6 maps each mechanism to what it controls and what it costs.

Table 7.6: Autograd Control Mechanisms: The recurring extension points the autograd engine exposes during graph traversal, what each controls, and the memory or correctness cost each carries.

Mechanism	What it controls	Cost or hazard
<code>retain_graph=True</code>	Keeps saved graph state available for another backward traversal	Retains memory that would otherwise be released; the amount depends on the graph and saved values
Gradient accumulation	Successive backward passes sum into <code>.grad</code> until <code>zero_grad()</code> resets them, enabling large effective batches	Forgetting the reset silently mixes gradients across optimization steps
Custom autograd.Function	User-defined forward and backward rules for operations outside the built-in library	Moves the differentiation contract (what to save, how to differentiate) to the implementer
Gradient hooks	Inspecting or modifying gradients mid-traversal (clipping, logging, debugging)	Runs arbitrary Python per registered tensor on every backward pass
<code>.detach()</code>	Cuts gradient flow at a chosen boundary (frozen layers, inference outputs)	The legacy <code>.data</code> attribute bypasses autograd and silently corrupts gradients; clone before in-place mutation

One mechanism deserves a closer look because it exposes the differentiation contract most directly. Custom autograd functions move part of that contract from the framework to the implementer: the

developer explicitly specifies what to save for the backward pass and how to compute gradients. Listing 7.14 shows the pattern.

Listing 7.14: Custom Autograd Function: Implement forward and backward methods to define custom differentiable operations, explicitly specifying tensors to save for gradient computation.

```
class MultiplyAdd(torch.autograd.Function):
    @staticmethod
    def forward(ctx, x, y, z):
        # Save tensors needed for backward
        ctx.save_for_backward(x, y)
        return x * y + z

    @staticmethod
    def backward(ctx, grad_output):
        # Retrieve saved tensors
        x, y = ctx.saved_tensors

        # Compute gradients using chain rule
        grad_x = grad_output * y # dL/dx = dL/dout * dout/dx
        grad_y = grad_output * x # dL/dy = dL/dout * dout/dy
        grad_z = grad_output # dL/dz = dL/dout * 1

        return grad_x, grad_y, grad_z

# Usage
x = torch.tensor([2.0], requires_grad=True)
y = torch.tensor([3.0], requires_grad=True)
z = torch.tensor([1.0], requires_grad=True)

output = MultiplyAdd.apply(x, y, z)
output.backward()

print(
    x.grad, y.grad, z.grad
) # tensor([3.]), tensor([2.]), tensor([1.]
```

These three principles connect directly to the framework’s role as a compiler for the silicon contract. The reverse-linked graph determines which operations the backward pass must execute (the O term). The memory-compute trade-off governs how much data the framework must move through the memory hierarchy (the D_{vol} term). The extensibility mechanisms, in turn, allow engineers to tune both terms for their specific workload. The interaction between autograd memory management and numerical precision leads naturally to mixed-precision training, which further reduces the D_{vol} term.

7.4.2.2 Mixed-precision training support

Mixed precision exploits a hardware asymmetry to improve two iron law terms simultaneously: Tensor Cores execute FP16 matrix multiplications at higher throughput than FP32 CUDA cores (raising R_{peak} and reducing the compute term $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$), while FP16 activations halve the memory footprint (reducing D_{vol}). Improving both terms simultaneously is rare; most optimizations improve one at the expense of the other.

Automatic mixed-precision APIs use reduced precision for eligible compute-intensive operations and higher precision where backend policy requires it. Matrix multiplications and convolutions commonly use FP16 on supported GPUs, while sensitive operations may use FP32; exact choices depend on the operation, device, and framework. This can preserve model quality and improve speed on suitable hardware. Because FP16 has a narrower dynamic range than FP32, gradients can underflow during backpropagation. Loss scaling addresses this by multiplying the loss before the backward pass, then dividing gradients by the same factor afterward.

Frameworks also support multiple precision modes including FP16, BF16,²¹ and TF32, NVIDIA’s

²¹ | **BF16 design rationale:** Developed by Google Brain circa 2018 specifically for TPU training stability, BF16 preserves FP32’s eight-bit exponent range while halving memory footprint—an explicit trade-off of mantissa precision (7 bits vs. FP16’s 10) for dynamic range. FP16’s smallest positive normal value is approximately 6.10×10^{-5} , but subnormal values extend to approximately 5.96×10^{-8} ; loss scaling is especially important on hardware or in modes that flush those subnormals to zero. BF16’s FP32-matched exponent largely avoids this class of gradient underflow, eliminating the need for loss scaling in most workloads, which is why BF16 and FP16 are not interchangeable: BF16 is preferred when training stability matters; FP16 is preferred when numerical precision matters more than gradient stability.

Tensor Core compute mode for FP32 matrix operations that keeps the FP32 exponent range while using lower mantissa precision. Each mode makes a different trade-off between range and precision. BF16 maintains FP32's dynamic range, simplifying training by eliminating most gradient underflow issues and often removing the need for loss scaling. Section 8.6.3 examines the mechanics of mixed-precision training in detail, including loss scaling algorithms, memory savings analysis, and numerical stability considerations. Listing 7.15 demonstrates PyTorch's mixed precision API: the autocast context manager automatically selects FP16 for compute-intensive operations while GradScaler prevents gradient underflow by dynamically scaling loss values.

Listing 7.15: Mixed-Precision API: Modern frameworks provide automatic mixed-precision support through context managers that handle precision selection and numerical stability.

```
import torch
from torch.amp import autocast, GradScaler

model = MyModel().cuda()
optimizer = torch.optim.Adam(model.parameters())
scaler = GradScaler("cuda")

for inputs, targets in dataloader:
    inputs, targets = inputs.cuda(), targets.cuda()
    optimizer.zero_grad()

    # Framework automatically selects precision per operation
    with autocast(device_type="cuda", dtype=torch.float16):
        outputs = model(inputs)
        loss = criterion(outputs, targets)

    # GradScaler handles gradient scaling for numerical stability
    scaler.scale(loss).backward()
    scaler.step(optimizer)
    scaler.update()
```

BF16 training typically does not require loss scaling. Comparing listing 7.16 with listing 7.15 line for line, the GradScaler construction and its scale, step, and update calls all disappear: BF16's FP32-matched exponent range reduces the gradient-underflow risk that often forces loss scaling, so the loop collapses back to an ordinary backward pass.

Listing 7.16: BF16 Training: BF16 maintains FP32's dynamic range, often avoiding the loss scaling used for FP16 training.

```
# BF16 training typically does not require loss scaling
with torch.autocast(device_type="cuda", dtype=torch.bfloat16):
    outputs = model(inputs)
    loss = criterion(outputs, targets)
loss.backward() # No GradScaler needed
optimizer.step()
```

Optimizer state and checkpointing Resuming training after interruption requires restoring model weights and optimizer state together: momentum buffers, adaptive learning rates, and gradient statistics. For Adam, optimizer state adds about $4\times$ the FP16 weight memory (two FP32 states per parameter), so weights plus optimizer state require about $5\times$ the FP16 weight footprint. A 7-billion-parameter model therefore requires approximately 70 GB total (14 GB weights + 56 GB optimizer state). Checkpoint size therefore bounds recovery speed after failure, connecting fault tolerance directly to the iron law's D_{vol} term.

Chapter 8 covers optimizer memory requirements and optimization strategies for large-scale training, where checkpoint size becomes a binding constraint. Frameworks provide the `state_dict()` interface to access optimizer state for serialization (listing 7.17), and resuming training requires loading both model parameters and optimizer state (listing 7.18).

Listing 7.17: State Dictionary Interface: Optimizers expose internal state through `state_dict()`, enabling serialization of momentum buffers and adaptive learning rate estimates for checkpointing.

```
import torch
import torch.nn as nn
import torch.optim as optim

model = nn.Linear(10, 5)
optimizer = optim.Adam(model.parameters(), lr=0.001)

# After training steps, optimizer accumulates state
loss = model(torch.randn(3, 10)).sum()
loss.backward()
optimizer.step()

# Access state for checkpointing
state = optimizer.state_dict()
# Contains: {'state': {...}, 'param_groups': [{'lr': 0.001, ...}]}
```

The mathematics of automatic differentiation were established decades before deep learning's resurgence. What changed was the systems engineering. Before framework automation, implementing gradient computation for a single fully connected layer meant writing separate forward and backward functions, manually tracking intermediate values, and verifying mathematical correctness across dozens of operations. A modern transformer involves hundreds of operations with complex dependencies; manual gradient derivation for attention, layer normalization, and residual connections would require months of careful work per architecture variant.

Listing 7.18: Checkpoint Save and Load: Save both model parameters and optimizer state to properly resume training with correct momentum and adaptive learning rate values.

```
# Saving checkpoint
checkpoint = {
    "epoch": epoch,
    "model_state_dict": model.state_dict(),
    "optimizer_state_dict": optimizer.state_dict(),
}
torch.save(checkpoint, "checkpoint.pt")

# Resuming training
checkpoint = torch.load("checkpoint.pt")
model.load_state_dict(checkpoint["model_state_dict"])
optimizer.load_state_dict(checkpoint["optimizer_state_dict"])
```

The breakthrough was turning this manual process into software infrastructure. A single matrix multiplication requires different gradient computations depending on which inputs require gradients, tensor shapes, hardware capabilities, and memory constraints. Autograd systems handle these variations transparently, which is why the rate of architectural innovation accelerated after frameworks matured. The mathematics did not change; software engineering made the mathematics practical to apply at scale.

7.4.2.3 Memory management in gradient computation

The memory strategies from section 7.4.1.2 (checkpointing, gradient accumulation) exist because reverse-mode differentiation requires preserving computational history. As listing 7.11 demonstrated, each layer adds an activation tensor that persists until the backward pass consumes it, creating a memory wave that peaks at the start of backpropagation and recedes as gradients are computed. Modern frameworks track the lifetime of each intermediate value automatically, freeing memory as soon as it is no longer needed. Even with precise lifetime tracking, however, a deeper problem remains: the cost of acquiring memory from the GPU in the first place.

The cost of raw GPU memory allocation provides a critical engineering lesson: production systems require memory abstraction. Device allocation can impose substantial overhead and synchronization. Modern frameworks therefore use **caching allocators** that retain and reuse device blocks rather than requesting memory for every tensor. Pooling reduces allocation calls and can mitigate fragmentation, but it cannot prevent fragmentation under all allocation patterns.

🔍 Systems Perspective 7.2: Caching allocator and utilization

Caching allocators primarily reduce allocation overhead and improve block reuse. Two considerations remain:

1. Allocation latency: Reusing a cached block avoids repeated device-allocation calls, whose cost and synchronization behavior depend on the allocator and runtime.
2. Fragmentation: Size classes and block splitting improve reuse but can leave free capacity in blocks that do not satisfy a request.

An out-of-memory error can reflect live allocations, allocator-reserved blocks, fragmentation, or another process. `nvidia-smi` reports device-level usage and is not sufficient by itself to diagnose the framework allocator; allocator-specific memory statistics are needed.

7.4.2.4 Production system integration challenges

A training iteration can run more slowly in production than in an isolated profile because the full step coordinates with the memory allocator, device manager, operation scheduler, data pipeline, and optimizer. The full-system cost includes work that an operator microbenchmark or isolated kernel profile does not measure. Gradient computation launches accelerator kernels and may require allocation, synchronization, or data movement according to the graph and implementation. These interactions can dominate small-operation workloads and remain measurable at scale. The gap between what the programmer writes (a five-line training loop) and what the system executes (allocations, kernel launches, and synchronization points) is the central tension of AD system design.

The severity of this overhead depends on a deeper architectural choice: how the AD system records and replays computation in the first place. A framework that builds a dynamic trace at runtime pays per-operation bookkeeping on every forward pass, and its caching allocator must accommodate allocation patterns that can vary with execution. A framework that captures a stable computation, by contrast, may pre-plan buffers and reuse storage across the forward and backward graphs because it can reason about lifetimes before execution. The gain depends on capture coverage, tensor lifetimes, shapes, and backend support.

🔍 Systems Perspective 7.3: Tape-based vs. transform-based autodiff

PyTorch (tape-based): Records operations on a dynamic “tape” during the forward pass. This is flexible and easy to debug but makes it hard for a compiler to see the whole graph at once for global optimization.

JAX (transform-based): Treats automatic differentiation as a high-level function transformation (`grad(f)`). JAX traces a pure function for abstract input signatures, allowing compatible transformations such as `jit(grad(f))` or `vmap(grad(f))` to compose before lowering. This exposes a larger transformed program to accelerator compilers without guaranteeing a single kernel or a universal speed advantage.

The AD system tracks dependencies so it can traverse independent graph branches in a valid order and accumulate gradients correctly. Independent host tasks or explicitly managed device streams can expose concurrency, but ordinary PyTorch autograd work on CUDA follows the framework’s stream semantics rather than automatically assigning every independent branch to a separate stream. Actual overlap depends on runtime scheduling, dependencies, and available device resources.

The memory and system integration challenges examined in section 7.4.1.3 (caching allocators, activation storage, and checkpoint overhead) affect all frameworks. Yet *how* frameworks implement

automatic differentiation in the first place varies significantly, with consequences for both optimization potential and developer experience. The distinction between tape-based and transform-based autodiff captures this architectural divergence. JAX²² exemplifies the transform-based approach, where composable function transformations replace imperative tape recording.

7.4.2.5 How different frameworks implement AD

The execution models covered in section 7.3, namely eager, static graph, and hybrid, directly shape how each framework implements automatic differentiation:

- PyTorch (Paszke et al. 2019) builds its autograd tape dynamically during forward execution, providing immediate debugging at the cost of graph-level optimization. The `grad_fn` chain mechanism detailed in section 7.4.2.1 enables flexible control flow but requires storing the complete graph until backward pass completion.
- TensorFlow (Abadi et al. 2016) (in its 1.x incarnation) performed symbolic differentiation during graph construction, enabling ahead-of-time optimization. Modern TensorFlow 2.x uses eager execution by default but provides `tf.function` for graph compilation when performance matters (TensorFlow Developers 2024a).
- JAX (Frostig et al. 2018) transforms functions rather than tracking operations. The `jax.grad()` transformation returns a new function that computes gradients, enabling composition with `jax.vmap()` for vectorization and `jax.jit()` for compilation. This approach requires pure functions but enables composable program transformations that chain differentiation, vectorization, and compilation in a single expression.

Autodiff implementation differences determine how much debugging visibility, compiler optimization, and deployment portability a team can expect from each framework. A recurring tension runs through every AD design decision: mathematical correctness demands storing computational history, but hardware imposes strict memory limits. Every framework resolves this tension differently, choosing which activations to checkpoint, which operations to fuse, and how aggressively to trade recomputation for memory. These choices determine which models can train on which hardware, making AD system design one of the most consequential engineering decisions in any framework.

Checkpoint 7.2: The systems cost of gradients

Training is inherently more expensive than inference because of Automatic Differentiation.

Computational reality

- ☐ **Reverse mode AD:** why is it efficient for a scalar loss with many parameters?
- ☐ **The activation tax:** can you explain how saved-value memory depends on graph depth, tensor shapes, and checkpointing?

Optimization mechanics

- ☐ **Gradient checkpointing:** how does recomputing activations save memory?

The execution and differentiation problems together enable the training loop: the execution model determines when computation happens, while automatic differentiation computes the gradients that drive learning. Both problems, however, quietly assume something that cannot be taken for granted: that the same code can run across diverse hardware. A model trained on an NVIDIA A100 must serve inference on a mobile phone's ARM CPU, a Google TPU, or a microcontroller with kilobytes of memory. The same `torch.matmul` call must dispatch to cuBLAS on one device and a hand-tuned ARM NEON kernel on another. This hardware diversity creates the third problem.

7.5 Abstraction Problem

This hardware diversity is architecturally fundamental. GPUs expose abundant throughput-oriented parallelism but have different memory and execution semantics from CPUs. TPUs favor regular tensor programs and compiler-visible shapes. A microcontroller has kilobytes where a server has

²² | **JAX:** JAX exposes `grad`, `jit`, and `vmap` as transformations on functions; they can be composed when their input and output contracts align, but order matters: vectorizing gradients is not generally the same operation as differentiating a vector-valued result. Compilation can produce fused regions, library calls, and multiple kernels rather than one universal kernel. Transformed functions should be pure; Python side effects may occur during tracing but are not represented in the compiled program, while supported callbacks express intentional runtime effects.

gigabytes. The **abstraction problem** therefore requires frameworks to hide this complexity behind a single programming interface while still enabling efficient utilization of each target’s unique capabilities.

The problem decomposes into two interacting dimensions. The first is *data representation*: how frameworks encode tensors, parameters, and computational state in forms that work across hardware. The second is *execution mapping*: how high-level operations translate into hardware-specific implementations. These dimensions are not independent concerns. The way data is represented (memory layout, precision, device placement) directly affects *what* execution strategies are possible. A tensor stored in row-major format on a GPU requires different kernels than one in column-major format on a CPU. A model quantized to INT8 enables entirely different execution paths than FP32.

Solving the abstraction problem requires sophisticated software infrastructure: tensor representations that encode both mathematical semantics and hardware constraints, intermediate representations that enable hardware-specific compilation, and runtime systems that manage data movement across the memory hierarchy. To make this concrete, trace what must happen when a programmer writes `model(input)`. The framework must resolve five decisions in rapid succession: data representation (tensor shape, memory layout, numeric precision), device placement (the bandwidth hierarchy connecting CPU, GPU, and accelerator memory), input delivery (data pipelines that sustain hundreds of MB/s to keep the accelerator fed), model organization (parameters, buffers, and submodules that must move together), and kernel execution (dispatch, scheduling, and resource optimization). These decisions build from the data container up to the hardware execution layer. Distributed placement previews the same abstraction boundary: frameworks expose placement scopes so system components can manage device boundaries, while chapter 8 analyzes the scaling algorithms and communication costs that make those scopes efficient.

7.5.1 Data structures and tensor abstractions

A ResNet-50 forward pass touches 25.6M parameters, produces intermediate activations at every layer, and must coordinate memory across CPU and GPU address spaces. Frameworks organize all of this data so that a single `model(input)` call executes millions of operations without the programmer managing a single pointer, by solving four problems in sequence: defining a universal data container (tensors), placing it on the right device (memory management), feeding data fast enough (data pipelines), and dispatching the right hardware kernel (core operations). The path runs from data representation to hardware execution.

Computational graphs specify the *logical* flow of operations, but data structures determine how those operations access and manipulate data in *physical* memory. This distinction matters because the same mathematical operation can differ by an order of magnitude in throughput depending on whether data is contiguous in cache, pinned for direct memory access (DMA) transfer, or scattered across pages.

The first step is the data container itself. Framework data structures must sustain memory bandwidth (hundreds of GB/s on modern GPUs), accommodate architectures from 1D sequences to 5D video tensors, and hide device management behind clean APIs. Tensors are the universal answer.

7.5.1.1 Tensors

At the foundation of every framework’s data representation lies a single abstraction: the tensor, an n -dimensional array that pairs numerical values with the information needed to interpret and place them.

Every computation in a neural network operates on tensors.²³ Training batches, activation maps, parameter gradients, and optimizer states are all tensors. This unified representation lets frameworks optimize a single data structure for hardware rather than managing separate containers for each role.



Definition 7.2: Tensor

Tensors are n -dimensional arrays with shape, data type, stride, and device metadata. Framework runtimes use this information to select compatible operations and hardware kernels.

²³ | **Tensor**: In mathematics, tensors obey coordinate-transformation laws; ML frameworks use the term more broadly for multidimensional arrays carrying shape, dtype, strides, device placement, and often automatic-differentiation metadata. A transpose commonly creates a strided view, while reshape may return either a view or a copy depending on the layout. These storage semantics make layout relevant to kernel selection and data movement.

1. **Significance:** A contiguous FP32 tensor of shape $[1024, 1024]$ occupies $1024 \times 1024 \times 4 = 4,194,304$ bytes (about 4.2 MB). A noncontiguous view may be accepted directly by a strided kernel or may require materialization when an operation demands a supported contiguous layout.
2. **Distinction:** NumPy arrays also carry shape, dtype, and stride metadata. Framework tensors additionally integrate device placement, automatic differentiation, and accelerator dispatch.
3. **Common pitfall:** Tensor operations may return a new allocation, a view sharing existing storage, or an in-place update, depending on the operation and API variant. Memory accounting must distinguish these cases.

The tensor abstraction consumes far more memory than model weights alone suggest. Engineers who estimate memory from parameter count alone allocate accordingly and encounter out-of-memory errors that seem inexplicable. A quick memory accounting makes the hidden cost concrete: gradients, optimizer momentum, and stored activations accompany every weight tensor.



Napkin Math 7.3: The administrative tax

The illustrative ResNet-50 breakdown in section 7.4.2.1 was $\sim 82.1\text{--}121.2\times$ larger than FP32 weight storage alone. The following scenario expands the accounting to gradients, optimizer state, and saved activations at billion-parameter scale.

Problem: How much memory does training state add beyond FP16 weights in this billion-parameter scenario?

Math:

1. Model weights: 2 GB.
2. Gradients: 2 GB (same size as weights).
3. Optimizer states (Adam): 8 GB ($4 \times$ FP16 weight memory for momentum and velocity stored in FP32).
4. Activations: For a batch size of 32 and a 100-layer network, reverse-mode AD retains the operation-dependent values required by derivative rules; activation checkpointing can omit selected values and recompute them during the backward pass.

$$\text{Activations} \approx B \times N_L \times S \times d_{\text{model}} \times n_{\text{saved}} \times 2 \text{ bytes}$$

For a 1024-token sequence, 1024-wide hidden state, and 1 saved tensor per layer: $32 \times 100 \times 1024 \times 1024 \times 1 \times 2 \approx \mathbf{6.7\text{GB}}$. Materialized attention terms can add a separate $B \times N_{\text{heads}} \times S^2$ component, which is why memory-efficient attention kernels matter.

Systems insight: A 2 GB model carries persistent training state before the first batch is processed (2 GB gradients + 8 GB optimizer state). During a training step, batch-dependent activations add another 6.7 GB, raising the peak “administrative tax” to ~ 16.7 GB beyond the weights. During training, data movement includes saving and retrieving these activations.

7.5.1.2 Tensor structure and dimensions

Tensors structure numerical data by adding organizational axes to linear memory layouts. Trace the rank expansion from left to right in figure 7.8, observing how rank 0 scalars evolve into rank 3 volume tensors.

In vision applications, raw image inputs map directly onto multi-dimensional tensor layouts. Examine the rank 3 tensor structure in figure 7.9, observing how red, green, and blue color channels stack along the channel dimension.

Framework tensors carry more than raw numbers. Each tensor stores **tensor metadata**, runtime information used to validate operations and select fast execution paths: a *shape* tuple (for example, $[64, 3, 224, 224]$ for a batch of images), a *dtype* (framework literals such as `float32`, `float16`,




Figure 7.8: Tensor Rank Hierarchy: Ascending tensor ranks (0 to 3) generalize multi-dimensional indexing from scalars to volume tensors. Each additional rank adds an organizational axis that the framework must linearize into physical memory addresses via strides.

or `int8`), and a `device` tag (CPU, cuda:0). A matrix multiplication, for instance, checks shape compatibility at dispatch time and uses the `dtype` to route to the correct hardware kernel, whether a standard FP32 GEMM or a Tensor Core FP16 path.




Figure 7.9: Image as an RGB Tensor: Decomposing a 2D color image into stacked height, width, and channel dimensions forms a rank-3 tensor. In ML systems, the ordering of these dimensions (such as channel-first vs. channel-last) determines memory access locality and cache efficiency during convolution operations.

Memory layout implementation introduces distinct challenges in tensor design. While tensors provide an abstraction of multi-dimensional data, physical computer memory remains linear. Stride patterns address this disparity by creating mappings between multi-dimensional tensor indices and linear memory addresses. These patterns significantly impact computational performance by determining memory access patterns during tensor operations. Figure 7.10 makes this concrete with a 2×3 tensor: follow the same six values as they map into two different linear orderings—row-major and column-major—and note how the stride values change to compensate.

Stride choices become performance choices. Row-major layout (used by NumPy, PyTorch) stores elements row by row, making row-wise operations more cache-friendly. Column-major layout (used by some BLAS libraries) stores elements column by column, optimizing column-wise access patterns. The stride values encode this layout information: in row-major layout for a 2×3 tensor, moving to the next row requires skipping three elements (`stride[0] = 3`), while moving to the next column requires skipping one element (`stride[1] = 1`).

These memory layout details have direct performance implications, but no ordering is universally superior. Efficient access depends on which dimension an operation traverses, the kernel's supported layouts, vectorization, and the target memory hierarchy. A mismatched layout may require strided access or a conversion, while a kernel designed for that layout can access it efficiently.

The `dtype` is the tensor-level lever that trades numerical range against data movement. The standard choice in machine learning has been FP32 precision, exposed in frameworks through `dtype` literals such as `float32`, offering a balance of precision and efficiency. Modern frameworks extend this with multiple numeric types for different needs. Integer types support indexing and




Figure 7.10: Tensor Memory Layout: A 2×3 tensor can be stored in linear memory using either row-major (C-style) or column-major (Fortran-style) ordering. Strides define the number of elements to skip in each dimension when moving through memory, enabling frameworks to calculate memory addresses for tensor[i,j] as $\text{base_address} + i \times \text{stride}[0] + j \times \text{stride}[1]$. The choice of memory layout significantly impacts cache performance and computational efficiency.

embedding operations. Reduced-precision types like FP16 enable efficient mobile deployment. INT8 precision allows fast inference on specialized hardware. The choice of numeric type affects both model behavior and computational efficiency: neural network training typically requires FP32 precision for critical accumulations to maintain stable gradient computations, while inference tasks can often use lower precision (framework dtype literals such as `int8` or even `int4`, corresponding to INT8 and INT4), reducing memory usage and increasing processing speed. Mixed-precision training approaches combine these benefits by using FP32 for critical accumulations while performing most computations at lower precision.

Type conversions become another point where abstraction meets physics. Operating on tensors with different types demands explicit conversion rules to preserve numerical correctness. These conversions introduce computational costs and risk precision loss. Frameworks provide type casting capabilities but rely on developers to maintain numerical precision across operations.

Tensors answer the data-representation problem by encoding shape, layout, and precision into a single abstraction. A perfectly shaped tensor on the wrong device, however, or one that must cross the PCIe-to-HBM bandwidth gap to reach the GPU, can erase every layout optimization. The next problem is placement: where data lives and how it moves.

7.5.1.3 Device and memory management


Tensors and their memory layouts establish what the framework computes with. Where that data physically resides, and how it moves between locations, determines whether computation happens at full speed or crawls.

Frameworks as the operating system interface While the high-level API focuses on math, the framework’s backend functions as the operating system of the Single-Machine Stack. It manages the two critical resources of a single node: compute scheduling and data movement.

The CUDA Runtime serves as this OS layer, providing the low-level primitives for launching kernels and managing device memory. The framework coordinates with this runtime to implement **DMA** over the PCIe bus. The bandwidth gap between the host (CPU) and device (GPU) is the primary “Data Loading Bottleneck”: section D.3.4 models this transfer, separating the two constraints, $T = L_{\text{lat}} + D_{\text{vol}}/\text{BW}$, that determine whether a given data-loading strategy is limited by per-transfer latency or by sustained bandwidth. Frameworks mitigate this through pinned memory (page-locked host memory), which enables DMA transfers without pageable-memory staging. This “HW/OS” interface is what makes high-throughput training loops possible on a single machine.

Every tensor resides on a specific device, and cross-device operations incur transfer costs that can dominate execution time. PCIe 4.0 delivers 32 GB/s between CPU and GPU, while HBM2e provides 2.04 TB/s within the GPU. This $63.7\times$ bandwidth gap means a single misplaced tensor transfer can erase the entire speedup from GPU acceleration.

Device placement matters for framework design because the framework must track where every tensor lives and enforce that operations only combine tensors on the same device. When data must move, the framework must decide whether to block execution or overlap the transfer with other work.



Interconnect bandwidth spans $\sim 64\times$: HBM far above NVLink, NVLink far above PCIe.

These decisions, invisible to most users, can move a training loop from transfer-bound execution toward substantially higher hardware utilization.

Three systems principles govern effective device and memory management: understanding the bandwidth hierarchy that constrains data movement, overlapping computation with communication to hide transfer latency, and using fine-grained synchronization to maintain correctness without sacrificing concurrency, supported by quantitative analysis grounded in the iron law’s data movement term. The bandwidth hierarchy provides the first constraint.

The cost of moving data between devices varies by orders of magnitude depending on the interconnect.²⁴ Table 7.7 shows transfer times for a 1000×1000 float32 tensor (4 MB)—roughly the size of a typical activation tensor in a moderately sized model. The numbers reveal why careless device placement can erase any speedup from GPU acceleration.

Table 7.7: Device Transfer Overhead: Transfer time for a 4 MB tensor across different interconnects. PCIe bandwidth shown is unidirectional (typical for GPU transfers), with full-duplex operation providing $2 \times$ total bandwidth. NVLink bandwidth is bidirectional (300 GB/s per direction). These transfer costs can dominate lightweight operations, making device placement critical.

Interconnect	Bandwidth	Transfer Time	Path
PCIe 3.0 x16	15.8 GB/s	0.254 ms	Host to device
PCIe 4.0 x16	32 GB/s	0.125 ms	Host to device
NVLink 3.0	300 GB/s per direction (600 GB/s bidirectional)	0.013 ms	GPU to GPU
GPU Memory	2039 GB/s	0.002 ms	On device

These numbers connect directly to the iron law of performance. Every cross-device transfer contributes to $(D_{\text{vol}}/\text{BW})$ at a fraction of on-device bandwidth. Dividing 1 GB by the stated PCIe 4.0 peak gives 31.2 ms, a bandwidth-only lower bound that omits protocol, launch, synchronization, topology, and contention overhead. Transfers can dominate lightweight or small-batch workloads, so deployed latency must be measured rather than inferred from peak bandwidth alone.

Every tensor should reside on the device where it will be consumed, and transfers should occur only when unavoidable. Frameworks track device placement for every tensor and raise errors when operations attempt to combine tensors from different devices, enforcing this discipline at the API level.

Overlapping computation and communication When transfers are unavoidable, the next optimization is to hide their latency by executing them concurrently with computation. Modern GPUs contain independent hardware units for computation (SM clusters) and data transfer (copy engines), enabling true simultaneous execution. The framework abstraction that exposes this hardware parallelism is the CUDA stream: an independent execution queue where operations execute sequentially within a stream but concurrently across streams.

Without explicit concurrency control, the GPU serializes all operations on a single default stream, leaving execution units idle while data transfers complete. By placing data transfers on one stream and computation on another, the effective latency approaches the theoretical minimum of $\max(\text{compute_time}, \text{transfer_time})$ rather than their sum. Stream-based overlap effectively hides the D_{vol}/BW penalty when computation is the longer operation (see listing 7.19).

The `non_blocking=True` flag requests an asynchronous copy with respect to the host when the backend supports it. For host-to-device copies, pinned memory is generally required to overlap the transfer with other device work; pageable memory may be staged internally, and the call’s host-blocking behavior is implementation-dependent. Correct overlap also requires compatible hardware, separate execution resources, and explicit dependency management.

The same synchronization pattern appears when model stages overlap across microbatches. Listing 7.20 shows each stage running on its own stream, with events enforcing only the producer-consumer dependencies needed for correctness.

²⁴ | **NVLink:** NVIDIA’s high-bandwidth GPU-to-GPU interconnect (see chapter 11), providing 600 GB/s bidirectional bandwidth (NVLink 3.0 on A100) compared to 64 GB/s for PCIe 4.0 x16. This $\sim 10 \times$ bandwidth advantage determines whether tensor parallelism, splitting one large tensor computation across multiple GPUs, is practical for a given model size: GPUs connected only by PCIe can make the D_{vol}/BW communication term dominate total training time, erasing the benefit of additional compute.



Overlapping copy and compute costs $\max(\text{copy}, \text{compute})$, not their sum.

Listing 7.19: Overlapping Computation and Transfer: Separate streams can overlap eligible copies and computation when dependencies are synchronized correctly. Pinned host memory generally enables host-to-device copies to overlap with other device work.

```
compute_stream = torch.cuda.Stream()
transfer_stream = torch.cuda.Stream()

# Transfer next batch while computing
# current batch
with torch.cuda.stream(transfer_stream):
    next_batch = next_batch_cpu.to("cuda", non_blocking=True)

with torch.cuda.stream(compute_stream):
    output = model(current_batch)
    loss = criterion(output, labels)

# Pinned host memory can enable transfer overlap
x_pinned = torch.randn(1000, 1000).pin_memory()
x_gpu = x_pinned.to("cuda", non_blocking=True) # Asynchronous

# A pageable source may require staging and
# should not be assumed to overlap
y_regular = torch.randn(1000, 1000)
y_gpu = y_regular.to("cuda", non_blocking=True)
```

Listing 7.20: Stage Overlap with Streams: Overlap multiple model stages across microbatches using streams and events for inter-stage synchronization.

```
# Pipeline parallelism: place stages on separate GPUs
devices = [torch.device(f"cuda:{i}") for i in range(3)]
stages = [
    Stage1().to(devices[0]),
    Stage2().to(devices[1]),
    Stage3().to(devices[2]),
]
streams = [torch.cuda.Stream(device=device) for device in devices]
events = [
    [torch.cuda.Event() for _ in range(num_microbatches)]
    for _ in stages
]
outputs = [[None] * num_microbatches for _ in stages]

for mb in range(num_microbatches):
    for stage_idx, (stage, stream, device) in enumerate(
        zip(stages, streams, devices)
    ):
        with torch.cuda.device(device), torch.cuda.stream(stream):
            if stage_idx > 0:
                # Wait for previous stage to complete
                # this microbatch
                events[stage_idx - 1][mb].wait()
                stage_input = outputs[stage_idx - 1][mb].to(
                    device, non_blocking=True
                )
            else:
                stage_input = inputs[mb].to(device, non_blocking=True)

            outputs[stage_idx][mb] = stage(stage_input)
            events[stage_idx][mb].record()
```

This overlap principle extends naturally to model-stage overlap within a single node. Different model stages on separate GPUs can process different microbatches concurrently, with each stage's computation overlapping the next stage's data reception (see listing 7.20). Chapter 8 later names and

analyzes the distributed-training strategies built from this scheduling pattern; here, the single-node implementation is enough to expose the synchronization principle that survives at larger scale. Once computation and communication overlap, the remaining challenge is ensuring correctness when operations complete out of order.

Synchronization and correctness Concurrent execution introduces ordering constraints. When one stream’s output becomes another stream’s input, the system must enforce a happens-before relationship without unnecessarily serializing independent work. Two synchronization mechanisms exist, with dramatically different performance implications.

Full device synchronization (`torch.cuda.synchronize()`) blocks all streams and the CPU until every queued operation completes. This creates a global serialization point that eliminates all overlap benefits. CUDA events provide the alternative: fine-grained synchronization that blocks only the dependent stream, allowing other streams and the CPU to continue execution (see listing 7.21).

Listing 7.21: CUDA Events for Synchronization: Events enable fine-grained producer-consumer patterns between streams without blocking the entire device.

```
# Create streams and event
stream1 = torch.cuda.Stream()
stream2 = torch.cuda.Stream()
event = torch.cuda.Event()

# Stream 1: producer
with torch.cuda.stream(stream1):
    result1 = expensive_computation(data1)
    event.record() # Mark completion point

# Stream 2: consumer (waits only for stream1's event)
with torch.cuda.stream(stream2):
    event.wait() # Block stream2 until event is recorded
    result2 = dependent_computation(result1) # Safe to use result1
```

The performance difference between these approaches is not incremental but categorical. Full synchronization after every operation converts a concurrent pipeline into a sequential one, entirely negating the hardware parallelism that streams expose. Event-based synchronization preserves the concurrent execution model while enforcing only the dependencies that correctness requires.

Device placement discipline protects the bandwidth hierarchy from accidental PCIe traffic. Every tensor carries a device attribute, and frameworks enforce a strict invariant: operations can only combine tensors on the same device. A `RuntimeError` results from mixing `cuda:0` and `cuda:1` tensors, preventing silent cross-device transfers.

The movement mechanism is explicit. `Tensor.to()` returns the original tensor when its dtype and device already match the request; otherwise it returns a converted copy unless `copy=True` forces a copy. `Module.to()` modifies registered parameters and buffers in place and returns the module.

The performance discipline follows from the bandwidth gap: allocate tensors on the target device from the start rather than creating them on CPU and transferring them, reuse GPU memory across iterations rather than reallocating it, and colocate inputs, labels, and model parameters on the same device to eliminate implicit transfers. At 32 GB/s, violating any of these principles inserts PCIe transfers into the critical path that can dominate a training iteration that otherwise runs at 2.04 TB/s on-device.

The same synchronization discipline has one operational trap: debug code often leaves `torch.cuda.synchronize()` in the hot path, turning an overlapped pipeline into a serialized one. When overlap remains poor, profiling must separate scheduling stalls from kernel inefficiency. NVIDIA Nsight Systems (`nsys profile`) shows CPU activity, GPU kernels, and memory transfers on one timeline. NVIDIA Nsight Compute (`ncu`) then explains kernel behavior with hardware counters.

Table 7.8 is the diagnostic map for that second step. SM means streaming multiprocessor, the GPU block that schedules groups of threads; a warp is one scheduled thread group.

Table 7.8: Nsight Compute Metrics: Key metrics for diagnosing ML kernels. Interpret each value against workload eligibility and kernel intent; Nsight Systems identifies which kernels dominate runtime, and Nsight Compute reveals why those kernels underperform.

Metric	Meaning	Optimization Target
SM Occupancy	Active warps/maximum warps	Increase parallelism if low
Memory Throughput	Achieved/peak bandwidth	Optimize memory access patterns
Compute Throughput	Achieved/peak FLOP/s	Reduce memory bottlenecks
Tensor Core Active	Time in Tensor Core ops	Verify mixed-precision utilization

7.5.1.4 Data pipelines and loading

Streams and events address placement and movement by overlapping transfers with computation so that the GPU rarely stalls on a single tensor. Scheduling alone, however, cannot help if data arrives too slowly in the first place. The next constraint is input delivery: data must arrive fast enough to sustain throughput. The core systems principle is straightforward: the data pipeline must sustain the accelerator's consumption rate. A GPU processing 1,000 images per second at 224 by 224 resolution requires approximately 150.5 MB/s of sustained raw uint8 image throughput. If the pipeline cannot maintain this rate, the accelerator idles and the effective utilization term in the iron law drops below 1.

Frameworks address this throughput requirement through three mechanisms. The first is *parallel worker processes*: the DataLoader spawns multiple CPU processes, each independently loading and preprocessing samples. When storage access, decoding, augmentation, or normalization makes one loader process too slow to feed the accelerator, multiple workers can overlap I/O wait with preprocessing. Once the input pipeline meets demand, extra workers add overhead rather than throughput. When `num_workers > 0`, the DataLoader distributes sample indices across workers through a shared queue, and workers push completed samples to a data queue that the main process assembles into batches.

The second mechanism is prefetching. With multiprocessing enabled, `prefetch_factor` controls how many batches each worker prepares in advance. Four workers with `prefetch_factor=2` can maintain up to eight prefetched batches, increasing the chance that input work overlaps accelerator computation. Prefetching cannot guarantee that the accelerator never stalls; storage, decoding, augmentation, and contention can still make production slower than consumption. The cost is additional host memory proportional to the queued data.

The third mechanism is *pinned memory for DMA transfers*. The `pin_memory=True` option places batches in page-locked host memory, which can avoid pageable-memory staging and enable host-to-device copies to overlap device work. For a batch of 64 FP32 images at $224 \times 224 \times 3$ (38.5 MB), dividing by the stated PCIe 4.0 x16 peak gives 1.2 ms; this is a bandwidth-only lower bound, not a measured pinned-transfer latency. The realized benefit over pageable memory depends on the platform, batch size, pipeline, and overlap. Pinned pages also reduce pageable system memory.

The DataLoader configuration is useful only when each parameter is tied to a bottleneck. In this configuration, `num_workers` enables parallel loading, `prefetch_factor` controls pipeline depth, and `pin_memory` enables DMA transfers. The worker count is a throughput/memory trade-off, not a universal constant. A practical starting point is setting `num_workers` equal to the number of available CPU cores, then adjusting based on whether loading is I/O-bound or CPU-bound. For I/O-bound workloads such as reading images from network storage, more workers overlap disk latency and improve throughput. For CPU-bound workloads involving heavy augmentation, the benefit saturates once all cores are in use. Too many workers waste memory, since each maintains a copy of the Dataset object.

Once throughput is high enough, worker process management becomes a correctness constraint. PyTorch assigns each worker a distinct PyTorch seed. A `worker_init_fn` is still useful when dataset or augmentation code also uses NumPy or Python's `random` module, as listing 7.22 demonstrates. Shared Python state is generally process-local, so modifications in one worker do not automatically propagate to the others or the main process; explicitly shared or memory-mapped storage is needed when workers must share mutable state.

- workers
- prefetch
- pinned

Each DataLoader knob relieves a specific input bottleneck.

The Dataset choice is another throughput decision because it determines how samples can be scheduled. PyTorch supports two dataset paradigms. *Map-style* datasets implement `__len__` and `__getitem__`, enabling random access to samples by index—this pattern works well for datasets that fit in memory or support efficient random access on disk. *Iterable-style* datasets implement `__iter__` instead, yielding samples sequentially for streaming data sources where random access is impractical. Map-style datasets support sampler-based shuffling, while iterable datasets must implement any ordering or buffered shuffling within the iterator or data source.

Listing 7.22: DataLoader Throughput Configuration: Each parameter addresses a specific throughput bottleneck. `num_workers` parallelizes I/O and preprocessing across CPU cores, `prefetch_factor` controls pipeline depth, and `pin_memory` enables DMA transfers to the GPU.

```
import random

import numpy as np
import torch
from torch.utils.data import DataLoader

def seed_worker(worker_id):
    worker_seed = torch.initial_seed() % 2**32
    np.random.seed(worker_seed)
    random.seed(worker_seed)

loader = DataLoader(
    dataset,
    batch_size=64,
    shuffle=True,
    num_workers=4, # Parallel worker processes (mechanism 1)
    prefetch_factor=2, # Batches prepared ahead per worker (mechanism 2)
    pin_memory=True, # Page-locked memory for DMA (mechanism 3)
    worker_init_fn=seed_worker, # Reproducible augmentation per worker
)

# Pipeline effect: while GPU processes batch N,
# 4 workers load batches N+1..N+8 into pinned memory,
# ready for DMA transfer when the GPU finishes.
```

Collation is the final place where representation choices affect throughput. The `collate_fn` parameter determines how individual samples are combined into batches. The default collation stacks tensors along a new batch dimension, which works when all samples have identical shapes. For variable-length data such as text sequences, custom collation handles padding, sorting by length, or creating attention masks—directly affecting both memory usage and training throughput.

DataLoaders, Datasets, and collation functions solve input delivery by sustaining accelerator-rate throughput through parallelism, prefetching, and DMA. These structures, however, handle only ephemeral data: samples flow through the pipeline once per epoch and are discarded. The next framework responsibility is persistent state, especially the model’s own weights when those weights exceed the memory of any single device.

Parameter structures

A GPT-3 scale model stores 175B parameters, occupying 350 GB in FP16. Managing these parameters across devices, keeping gradients synchronized, and maintaining optimizer state (Adam state alone can add about 4× the FP16 weight memory, as the Administrative Tax notebook showed) is a core framework responsibility. Because parameters persist throughout training and inference, frameworks organize them into compact structures that minimize memory while enabling fast read and write access. During multi-GPU training, frameworks may replicate parameters across devices for parallel computation while keeping a synchronized master copy; parameter-server systems are one communication-efficient design for workers to read and write globally shared parameters (Li et al. 2014). Synchronizing multi-billion parameter models can require transferring tens of GB of

gradients per step, which is why frameworks expose communication backends that can synchronize tensors efficiently. Chapter 8 later names the specific collective operations and scaling strategies.

Parameter structures must also adapt to varying precision requirements. Training typically uses FP32 for gradient stability, but inference and large-scale training increasingly use FP16 or INT8. Frameworks implement type casting and mixed-precision management to enable these optimizations without compromising numerical accuracy.

Distributed execution contexts

The computational graph defines *what* to compute, but *where* and *how* that computation runs across devices is the job of execution contexts. On a single node, execution contexts manage CUDA streams and events (introduced in section 7.5.1.3) to overlap computation and data transfer across GPUs.

When training scales beyond a single machine, these same abstractions extend to named groups of devices. Frameworks use constructs like `ProcessGroup` (PyTorch) or `Mesh` (JAX) to describe which tensors and operations belong together, relying on communication libraries like NCCL to execute collective primitives—such as **ring-AllReduce** for gradient aggregation and **AllGather** for parameter gathering. The runtime manages these primitives to maximize interconnect utilization: intra-node NVLink links deliver up to 900 GB/s of bi-directional bandwidth per GPU, whereas inter-node InfiniBand or Ethernet networks drop to 50–400 Gbps (6.25–50 GB/s), making collective communication efficiency the primary bottleneck in distributed training scaling. The important framework idea is the abstraction boundary: user code names placement relationships, and the framework preserves those relationships while the hardware path changes underneath.

These concepts appear here because they shape framework API design even before the book asks the reader to reason about distributed-training algorithms. The details of gradient synchronization, communication topologies, and fault tolerance build on these foundations later. For now, the only point needed is placement expressiveness: when models exceed single-device memory, frameworks must give the training system more than one way to place work. A GPT-3 scale model, for instance, cannot fit on a single GPU—its 175B parameters alone require 350 GB in FP16, far exceeding any GPU’s memory. Figure 7.11 previews the framework placement idea without requiring the full distributed-training machinery yet: a system can split work across layer groups, across replicated batches, or within very large tensors. Section 8.7 later formalizes these dimensions as training strategies and analyzes their communication costs.




Figure 7.11: Device Placement Dimensions: A grid of eight accelerator clusters arranged in two rows and four columns, each containing stacked computational units. The braces preview three ways a framework can place work when a model no longer fits one device: across layer groups, across replicated batches, and within very large tensors.

The data structures examined so far—tensors, device managers, data pipelines, parameter structures, and distributed execution contexts—define *what* data a framework manages and *where* it lives. The remaining abstraction problem is execution: what actually runs on the hardware.

7.5.2 Core operations

When an engineer writes `y = torch.matmul(x, w)`, the gap between Python and the GPU is larger than it appears. The gap between a single line of Python and thousands of parallel GPU threads is bridged by three groups of operations working in coordination. Figure 7.12 shows the full stack:

hardware abstraction operations manage platform-specific execution, basic numerical operations implement mathematical computation, and system-level operations coordinate scheduling, memory, and resources across the graph. The prose follows that build-up from the hardware-specific kernel path toward system orchestration.




Figure 7.12: Core Operations Stack: Three functional tiers bridge high-level Python code to physical execution. Hardware operations manage low-level kernel dispatch and memory transfers, basic numerical operations execute dense GEMM and element-wise routines, and system-level operations coordinate graph scheduling and memory lifetimes across accelerators.

7.5.2.1 Hardware abstraction operations

The hardware abstraction layer isolates framework code from platform-specific details. It solves three concrete problems: selecting the right compute kernel, moving data through the memory hierarchy, and coordinating execution across processing units.

Compute kernel management The kernel manager dispatches each supported operation to an implementation for the current backend. For matrix multiplication, this may be a CPU BLAS backend such as Intel oneMKL or OpenBLAS (Intel Corporation 2026a; OpenBLAS Project 2026), cuBLAS on NVIDIA GPUs (NVIDIA 2024a), or accelerator-specific tensor instructions. Selection depends on dimensions, dtype, layout, backend, and library heuristics. A 4096×4096 FP16 GEMM on an A100 may use a Tensor Core path whose published peak is represented here by 312 TFLOP/s (Choquette et al. 2021; NVIDIA Corporation 2020a). If the backend has no implementation for an operator, execution may fall back, be partitioned, or fail, depending on the framework and runtime.

Memory system abstraction The memory abstraction layer moves tensors among pageable or pinned host memory, device memory, and unified memory, and transforms data layouts to match hardware preferences. A convolutional layer, for example, may store activations in NCHW format (batch, channels, height, width) on NVIDIA GPUs but convert to NHWC for Apple’s Metal backend. Alignment requirements vary from 4 bytes on CPUs to 128 bytes on some accelerators, and misaligned access can halve effective memory bandwidth. The runtime also preserves dependency ordering when multiple execution units access the same tensor, preventing races during concurrent operations.

Execution control The execution controller coordinates work across multiple processing units and memory spaces. On a modern GPU, graph runtimes or explicit stream use can overlap independent kernels when dependency analysis proves they are ready and the kernels leave enough resources unused to run concurrently. Eager default-stream execution often serializes this work instead. The controller inserts synchronization barriers where data dependencies require them, tracks event completions to trigger dependent operations, and routes hardware errors (ECC failures, timeout watchdogs) to the framework’s error handling path.

7.5.2.2 Basic numerical operations

With hardware abstraction managing the platform-specific details, frameworks build a layer of mathematical operations on top. GEMM dominates ML computation. Section C.1.4 derives how GEMM arithmetic intensity scales with matrix dimension, predicting whether a given layer is compute bound or memory bound before any profiler runs. The operation $C = \alpha AW + \beta C$ accounts for the vast majority of arithmetic in neural networks: a single ResNet-50 forward pass performs approximately 8.2 GFLOP, nearly all of which reduce to GEMM. Frameworks optimize GEMM through cache-aware **tiling** (splitting matrices into blocks that fit in L1/L2 cache), loop unrolling for instruction-level parallelism, and shape-specific kernels. Fully connected layers use standard

dense GEMM, while convolutional layers use im2col transformations that reshape input patches into matrix columns, converting convolution into GEMM.

Beyond GEMM, frameworks implement BLAS operations (AXPY for vector addition, GEMV for matrix-vector products) and element-wise operations (activation functions, normalization). Element-wise operations are *individually cheap* but *collectively expensive* due to memory bandwidth. Each operation reads and writes the full tensor, so a sequence of five element-wise operations on a 100 MB tensor moves 1 GB of data. Fusing those five operations into a single kernel reduces memory traffic to 200 MB, a $5\times$ bandwidth savings that directly translates to faster execution.

Numerical precision adds another dimension. Training in FP32 uses 4 bytes per parameter; quantizing to INT8 reduces this to 1 byte, cutting memory by $4\times$ and enabling $2\text{--}4\times$ throughput improvements on hardware with INT8 acceleration. Training typically keeps numerically sensitive accumulations in higher precision, while inference can often run many operations in FP16 or INT8 with little quality loss. Frameworks maintain separate kernel implementations for each precision format and handle workflows where different layers operate at different bit widths within a single forward pass.

7.5.2.3 System-level operations

Hardware abstraction and numerical operations provide the building blocks; system-level operations orchestrate them. The system layer ties scheduling, memory management, and resource optimization into a coherent execution engine.

The operation scheduler uses graph dependencies to identify legal execution orders and possible concurrency. Static capture can expose more of the dependency structure before execution, while eager runtimes discover work incrementally. Neither visibility guarantees an optimal schedule or simultaneous execution: overlap depends on streams, compiler and runtime decisions, resource availability, and the operations themselves.

The memory manager allocates and reclaims GPU memory across the computational graph's lifetime. Model parameters (a 7-billion-parameter model consumes approximately 14 GB in FP16) persist for the entire training run, while activation tensors live only until the backward pass consumes them. PyTorch's caching allocator maintains a memory pool, subdividing and reusing freed blocks without returning them to CUDA, which avoids repeated `cudaMalloc` calls that can cost tens of microseconds and may be worse when they synchronize the device. For models that exceed GPU memory, the manager can apply checkpointing by discarding selected activations during the forward pass and recomputing them during the backward pass. The policy question—which activations to keep, and how much recomputation to tolerate—belongs to the training pipeline; the framework abstraction is what makes that policy executable.



Checkpoint 7.3: Hardware abstraction

The abstraction problem is the bridge between *portable code* and *efficient execution*.

- ☐ **Two dimensions:** can you distinguish data representation (layout, dtype, placement) from execution mapping (kernel selection, scheduling), and explain how they constrain each other?
- ☐ **Kernel dispatch:** can you explain why the same high-level operation (for example, GEMM) needs multiple implementations (CPU vector path vs. GPU Tensor Core path) and how shapes/dtypes affect the choice?
- ☐ **Memory abstraction:** can you explain why frameworks use caching allocators and layout transforms (NCHW $\leftrightarrow$ NHWC) rather than calling device allocators on every tensor?
- ☐ **Execution control:** can you describe what the runtime is doing when it overlaps independent work (streams) and inserts synchronization only where dependencies require it?

The resource optimizer integrates scheduling and memory decisions with backend kernel selection. Matrix-multiplication libraries choose among tiled GEMM implementations according to shapes, dtype, layout, workspace, and hardware. Winograd is instead a convolution-specific transform, and

Strassen is not a routine alternative selected by mainstream ML GEMM dispatch. A poor schedule can leave resources idle, while memory-pool fragmentation or excessive live state can cause an out-of-memory error even when aggregate device capacity appears sufficient.

The preceding sections examined what happens beneath the API surface: tensors manage data layout, streams overlap computation with communication, and kernel dispatch routes operations to hardware. These mechanisms operate at the level of individual tensors and operations—the raw materials of machine learning computation. Practitioners, however, rarely write code at this level. A ResNet-50 has 25.6M parameters organized into dozens of layers; manually tracking each tensor, registering it with an optimizer, and handling device placement would be error-prone and tedious. The abstraction problem is not fully solved by hardware-level mechanisms alone; it also requires a programming model that organizes these low-level primitives into the clean APIs that practitioners actually use.

Individual operations—matrix multiplications, activations, normalizations—are the atoms of deep learning computation. Building models from individual operations, however, would be like building a house from individual atoms. Frameworks need an organizational abstraction that lets engineers compose operations into reusable, nestable building blocks. That abstraction is the *module*.

7.6 nn.Module Abstraction

The hardware-facing half of the abstraction problem—tensors, kernels, streams, and memory managers—makes individual operations fast on diverse silicon. A ResNet-50, however, contains fifty layers, each with multiple parameter tensors, buffers, and mode-dependent behaviors. Manually wiring each tensor to the correct device, registering it with an optimizer, toggling dropout behavior between training and inference, and serializing state for checkpointing across every layer would drown practitioners in bookkeeping that has nothing to do with model design. The upper layer of the abstraction problem is organizational: composing thousands of low-level primitives into the clean, composable APIs that practitioners actually use.

Every major framework answers this question through a *module abstraction* that bundles parameters, forward computation, and state management into a single reusable unit. PyTorch’s `nn.Module`²⁵ provides an instructive case study because its design patterns recur across frameworks: Keras uses similar layer abstractions (Chollet 2018), JAX’s Flax employs analogous module structures, and TensorFlow’s functional API shares conceptual parallels. Three enduring design principles recur regardless of syntax or programming paradigm.

7.6.1 Automatic parameter discovery

A modern neural network may contain millions of trainable parameters spread across dozens of layers. Without automation, a programmer would need to enumerate every parameter tensor and pass it to the optimizer manually, an error-prone process that scales poorly with model complexity. Frameworks solve this through **automatic parameter discovery**: the system walks the module tree, collecting every parameter tensor so the optimizer can update them in a single call.

This is a tree traversal problem at its core. PyTorch’s `Module.__setattr__` registers an assigned `nn.Parameter` or child `nn.Module` in the module’s internal mappings. A call to `.parameters()` recursively traverses registered submodules and yields registered parameters. Other frameworks expose analogous parameter trees or collections, though their mechanisms differ.

The systems consequence is significant. Automatic parameter discovery gives the optimizer a complete parameter set or parameter groups, enabling grouped, foreach, or fused update paths when the framework and backend support them. The immediate win is correctness and composability: newly attached submodules participate in device moves, serialization, and optimizer construction without a separate manual registry. It also lets optimized update paths operate over parameter collections without per-parameter user code. Listing 7.23 demonstrates the core mechanism: attribute assignment triggers registration, and `.parameters()` returns all discovered tensors.

The distinction between *parameters* and *buffers* illustrates a subtlety of discovery. Registered parameters appear in `.parameters()` whether trainable or frozen; `requires_grad` separately controls whether autograd accumulates a gradient for them. Registered buffers move with the module and can appear in its state dictionary but are excluded from `.parameters()` and ordinary optimizer discovery. Batch-normalization running statistics are a common buffer use case.

²⁵ | **nn.Module**: The “design patterns recur” claim holds because `nn.Module` solves a universal organizational problem: it automatically registers any assigned submodule or parameter into a hierarchical tree, enabling a single `.to('cuda')` call to recursively place millions of parameters onto a GPU. Keras layers, JAX Flax modules, and TensorFlow’s `tf.Module` all implement the same tree-walking pattern. Without it, managing model state would require manual bookkeeping that scales linearly with architectural depth, a cost that grows prohibitive for models with hundreds of layers.

Listing 7.23: Parameter Registration: Automatic parameter tracking through attribute assignment enables optimizer access to all trainable weights without manual enumeration.

```
import torch
import torch.nn as nn

class CustomLayer(nn.Module):
    def __init__(self, input_size, output_size):
        super().__init__()
        self.weight = nn.Parameter(
            torch.randn(output_size, input_size)
        )
        self.bias = nn.Parameter(torch.randn(output_size))
        self.register_buffer("running_mean", torch.zeros(output_size))

    def forward(self, x):
        return torch.matmul(x, self.weight.t()) + self.bias

layer = CustomLayer(10, 20)
# Framework discovers both parameters automatically:
for name, param in layer.named_parameters():
    print(f"{name}: shape {param.shape}")
```

Table 7.9 shows how the same principle manifests differently across frameworks. Despite syntactic differences, all frameworks solve the same problem: enabling optimizers to discover and update trainable parameters while preserving nontrainable state across forward passes.

Table 7.9: Parameter Discovery APIs Across Frameworks: Each framework exposes a different surface for distinguishing trainable parameters from nontrainable buffers, but all four solve the same underlying problem of letting optimizers find weights to update while preserving state (running statistics, frozen tensors) across forward passes. Across all four rows, the durable pattern consists of a trainable-parameter accessor for the optimizer and a separate nontrainable-state channel for values that must persist but must not receive gradients.

Framework	Parameter Access	Nontrainable State
PyTorch	model.parameters()	register_buffer()
Keras	layer.trainable_weights	layer.non_trainable_weights
JAX/Flax	variables["params"] after variables = model.init(key, x)	Separate variable collections (for example, batch_stats)
TensorFlow	module.trainable_variables	module.non_trainable_variables

Mode-dependent behavior

Training and inference require different computational behavior from the same model graph. During training, dropout layers randomly zero elements with probability $p_{\text{drop}} = \text{Pr}(\text{drop})$ to regularize the network, while during inference those same layers must perform identity mapping to produce deterministic outputs. Batch normalization uses per-batch statistics during training but switches to accumulated running statistics during inference. If these behavioral changes are left to the programmer, forgetting a single mode switch produces silently incorrect predictions in production. Frameworks solve this with a *state flag* that propagates through the module hierarchy. A single call to `.eval()` on the root module recursively sets `self.training = False` on every descendant, and each layer queries this flag to select its behavior. This is an instance of a broader systems principle: the same computation graph must produce different execution behavior depending on context. Compilers face the same challenge when the same source code must produce debug builds (with bounds checking and symbol tables) vs. release builds (with aggressive optimization). The flag-propagation pattern ensures correctness by centralizing the mode decision at the root rather than requiring per-layer coordination.

This principle extends to parameter freezing for transfer learning. Setting `requires_grad=False` prevents autograd from accumulating gradients in those parameters and removes their gradient storage. It does not necessarily remove the layer’s backward computation because gradients may still need to pass through its operations to reach trainable inputs or earlier parameters. Savings therefore depend on where the frozen region lies and which tensors require gradients.

7.6.2 Hierarchical composition and serialization

Complex models compose from reusable submodules, creating a tree structure. A ResNet is not implemented as a monolithic block of operations but as a hierarchy: the root module contains a sequence of residual blocks, each block contains convolution layers and normalization layers, and each layer contains parameter tensors. This hierarchical composition must support two critical operations made concrete in listing 7.24: recursive parameter collection for training and **state serialization** for checkpointing and deployment.

Listing 7.24: Nested Module Composition: Hierarchical module composition enables recursive parameter collection and a flat state mapping across the module tree.

```
import torch
import torch.nn as nn

class ResidualBlock(nn.Module):
    def __init__(self, channels):
        super().__init__()
        self.conv1 = nn.Conv2d(channels, channels, 3, padding=1)
        self.bn1 = nn.BatchNorm2d(channels)
        self.conv2 = nn.Conv2d(channels, channels, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(channels)

    def forward(self, x):
        residual = x
        x = torch.relu(self.bn1(self.conv1(x)))
        x = self.bn2(self.conv2(x))
        return torch.relu(x + residual)

class ResNet(nn.Module):
    def __init__(self, num_blocks, channels=64):
        super().__init__()
        self.conv_in = nn.Conv2d(3, channels, 7, padding=3)
        self.blocks = nn.ModuleList(
            [ResidualBlock(channels) for _ in range(num_blocks)]
        )
        self.fc = nn.Linear(channels, 10)

    def forward(self, x):
        x = self.conv_in(x)
        for block in self.blocks:
            x = block(x)
        x = x.mean(dim=[2, 3]) # Global average pooling
        return self.fc(x)

model = ResNet(num_blocks=4)
total = sum(p.numel() for p in model.parameters())
print(f"Total parameters: {total}")
# state_dict() flattens the tree: 'blocks.0.conv1.weight', etc.
print(list(model.state_dict().keys())[:4])
```

Hierarchical composition mirrors the hardware memory hierarchy in a systems-relevant way: each submodule’s parameters can be loaded independently, enabling placement across devices. When

a model is too large for a single GPU, the framework can assign different subtrees of the module hierarchy to different devices, with the tree structure providing natural partition boundaries.

The `state_dict()` method produces an ordered mapping whose dotted keys (for example, `blocks.0.conv1.weight`) encode locations in the module hierarchy. A checkpoint for 7 billion FP16 parameters contains approximately 14 GB of weight payload before metadata or additional state, but the file format and storage path determine how those bytes are serialized. `load_state_dict()` copies matching parameter and buffer values into an already-constructed module and reports missing or unexpected keys according to its strictness setting; it does not reconstruct the module hierarchy. Cross-framework exchange requires a compatible graph and operator representation in addition to weights.

The hierarchical structure also enables module-level traversal for systematic operations. Methods like `.named_modules()` iterate the entire tree, supporting bulk transformations such as replacing all BatchNorm layers with GroupNorm or applying Xavier initialization to every Linear layer. These traversal operations depend on the same tree structure that enables parameter discovery, illustrating how a single design decision propagates benefits across multiple use cases.

These three principles, automatic parameter discovery, mode-dependent behavior, and hierarchical composition with serialization, are not PyTorch-specific. Every framework must solve them. Keras layers, JAX’s Flax modules, and even functional approaches all address the same problems of parameter management, state tracking, and compositional design. The differences lie not in *what* problems they solve but in *how* they prioritize among competing solutions. Two practical patterns show how the principles become system controls: selective parameter freezing reduces unnecessary gradient work for transfer learning (listing 7.25), and module hooks provide noninvasive inspection (listing 7.26).

Listing 7.25: Parameter Freezing: Demonstrates selective parameter freezing for transfer learning, where pretrained layers remain fixed while new layers train.

```
from torchvision.models import ResNet18_Weights, resnet18

# Freeze all parameters in a pretrained
# model
pretrained_model = resnet18(weights=ResNet18_Weights.DEFAULT)

for param in pretrained_model.parameters():
    param.requires_grad = False

# Replace final layer with trainable parameters
pretrained_model.fc = nn.Linear(512, 10) # New layer is trainable

# Only fc.parameters() will receive
# gradients during training
optimizer = torch.optim.Adam(
    filter(lambda p: p.requires_grad, pretrained_model.parameters()),
    lr=0.001,
)
```

Module hooks are the inspection counterpart to parameter freezing: they intercept intermediate computations without modifying model code, enabling gradient flow diagnosis and activation monitoring. Listing 7.26 illustrates both hook types.

Together, these patterns—parameter discovery, freezing, and hooks—demonstrate how the three principles translate into practical APIs. These `nn.Module` patterns illustrate PyTorch’s approach to the abstraction problem. PyTorch, however, is only one of several major frameworks, and its choices (mutable state, class inheritance, eager execution by default) are not the only valid design points. TensorFlow centralizes state differently, and JAX avoids mutable state entirely. These are not superficial API differences; they reflect deeply different answers to the chapter’s three opening problems.

Listing 7.26: Module Hooks: Shows forward and backward hooks for inspecting activations and gradients during training.

```

import torch
import torch.nn as nn

model = nn.Sequential(nn.Linear(10, 20), nn.ReLU(), nn.Linear(20, 5))

# Forward hook to inspect activations
def forward_hook(module, input, output):
    print(
        f"Layer: {module.__class__.__name__}, "
        f"Output shape: {output.shape}, "
        f"mean={output.mean():.3f}, "
        f"std={output.std():.3f}"
    )

# Backward hook to inspect gradients
def backward_hook(module, grad_input, grad_output):
    print(f"Gradient norm: {grad_output[0].norm():.3f}")

# Register hooks on specific layer
handle_fwd = model[0].register_forward_hook(forward_hook)
handle_bwd = model[0].register_full_backward_hook(backward_hook)

# Execute forward and backward pass
x = torch.randn(32, 10)
y = model(x)
loss = y.sum()
loss.backward()

# Remove hooks when done
handle_fwd.remove()
handle_bwd.remove()

```

7.7 Framework Platform Analysis

A team that prototypes quickly but cannot deploy the resulting model, or deploys reliably but cannot debug training failures, has run into a framework design trade-off rather than a missing API call. Each major framework solves all three core problems but emphasizes a different path through them. TensorFlow couples eager development to graph capture and a broad deployment ecosystem; PyTorch begins with eager execution and adds capture, compilation, and export paths; JAX organizes differentiation, vectorization, and compilation as composable function transformations. These are architectural emphases, not exclusive capabilities, and their value depends on the workload and deployment target.

7.7.1 TensorFlow: Eager development with graph deployment

TensorFlow's architecture spans eager development and graph-based optimization across hardware from cloud TPUs to microcontrollers. TensorFlow 2 executes eagerly by default. Performance-sensitive functions can be captured with `tf.function`, and deployment formats route captured computation to server, mobile, and browser runtimes. This dual path preserves the production graph capabilities inherited from TensorFlow 1.x without requiring every TensorFlow 2 program to build a complete graph before it runs.

Within a captured `tf.function`, TensorFlow can apply graph optimizations such as constant folding, operator fusion, and layout planning. Figure 7.13 maps the broader training-to-deployment pipeline from data preprocessing and distributed training to SavedModel export, TensorFlow Serving, TensorFlow Lite, TensorFlow.js, and language bindings.

While TensorFlow 2.0 introduced eager execution to bridge the gap between research and production, TensorFlow 2.x still exposes `tf.function` as the graph-conversion path for performance-

● TensorFlow

● PyTorch

● JAX

Each framework optimizes for a different system bottleneck.




Figure 7.13: TensorFlow Training-to-Deployment Pipeline: The SavedModel format serves as the central decoupling boundary between training and deployment. It captures graph structure, weights, and signatures from heterogeneous training setups (left) and exports them directly to target-specific runtimes (right)—including TF Serving, TF Lite, and TF.js—without requiring Python in production.

sensitive code (TensorFlow Developers 2024a). Its core strength remains the robust, compiled path from research to global-scale deployment. Chapter 8 and chapter 13 later examine the scaling and production infrastructure that use these export paths.

7.7.2 PyTorch: The eager research standard

Where TensorFlow’s graph-first approach prioritizes *production optimization*, PyTorch makes the opposite trade-off: it prioritizes *developer experience*. PyTorch’s architecture represents a sharply different answer to the execution problem, built on dynamic graphs (or “Define-by-Run”). Instead of building a blueprint before execution, PyTorch builds the computational graph on-the-fly as the code runs. Facebook AI Research (FAIR) adopted this design because researchers need immediate feedback when experimenting with novel architectures; the define-then-run cycle of static graphs introduced a compilation delay that slowed the rapid prototyping essential to research workflows.

PyTorch’s approach fits exploratory research for the same reason: it treats deep learning as standard Python programming. Developers can use Python loops, conditionals, and debuggers (like `pdb`) directly within a model’s forward pass, with no special syntax, no separate compilation step, and no waiting to see if the code works. Eager execution enables rapid iteration and intuitive model design, which is essential when architectures and training objectives are still changing.

PyTorch’s answer to the differentiation problem is tape-based autograd (section 7.4.2.1): flexible and debuggable, but harder to optimize globally because the tape is rebuilt each iteration. Its answer to the abstraction problem is more pragmatic than comprehensive: strong GPU support through cuBLAS and cuDNN, with deployment paths including `torch.export`, edge-targeted ExecuTorch, Open Neural Network Exchange (ONNX), and specialized runtimes.

The trade-off is therefore a more fragmented deployment path. Because the graph is dynamic, the framework cannot easily perform global optimizations before execution. A model that works perfectly in development may hit performance walls in production when dispatch overhead dominates small operations. To bridge this research-to-production gap, PyTorch introduced graph-capture and compilation paths, from TorchScript historically to `torch.compile` and export workflows, which allow developers to capture a dynamic model and turn it into an optimized representation for deployment. This evolution shows how an eager framework can move toward the production end of the compilation continuum while preserving the interactive experience that motivated the design.

7.7.3 JAX: The functional transformation engine

PyTorch’s eager execution and TensorFlow’s graph compilation represent two points on a spectrum, yet both share an imperative programming heritage where computation proceeds as a sequence of stateful operations. JAX represents a radically different user-facing approach, one built on functional

programming principles and composable program transformations rather than object-level tapes or user-authored graph APIs (Bradbury et al. 2018). Developed by Google Research, JAX is especially useful for work requiring custom differentiation, advanced optimization research, and large-scale distributed training.

JAX’s architecture treats differentiation, vectorization, and compilation as transformations on functions. The `jax.grad` function returns a function that computes gradients, which can then be transformed again when the interfaces are compatible. For example, `vmap(grad(f))` can compute per-example gradients and `jit(vmap(grad(f)))` can compile that batched computation. Transformation order is semantically meaningful rather than arbitrary.

JAX’s functional paradigm shifts the programming model from “tracking state through objects” to “transforming pure functions.” While PyTorch and TensorFlow expose autograd primarily through dynamic tapes or graph-compilation paths, JAX asks users to write pure Python functions and then applies transformations to those functions. Automatic differentiation, vectorization, and JIT compilation are all *program transformations* that can compose. Listing 7.27 demonstrates this approach.

Listing 7.27: JAX Function Transformations: JAX treats differentiation, vectorization, and compilation as composable function transformations rather than user-authored graph operations.

```
import jax
import jax.numpy as jnp

def loss_fn(params, x, y):
    pred = jnp.dot(x, params["w"]) + params["b"]
    return jnp.mean((pred - y) ** 2)

# Transform: compute gradients
grad_fn = jax.grad(loss_fn)

# Transform: vectorize over batch dimension
batched_grad = jax.vmap(grad_fn, in_axes=(None, 0, 0))

# Transform: compile to XLA
fast_batched_grad = jax.jit(batched_grad)

# Compose all three: fast, batched gradient computation
```

This functional approach requires **pure functions** (no side effects) and **immutable data** (arrays cannot be modified in place). These constraints may seem restrictive coming from PyTorch’s mutable object model, but they enable formal guarantees: the compiler can safely reorder, fuse, and parallelize operations because function outputs depend only on inputs. The restriction is the feature; purity is what makes transformation composition possible.

JAX’s power emerges from composition. `jax.grad` returns a gradient function; `jax.vmap` can vectorize a compatible function over mapped axes; and `jax.jit` can trace and compile a stable transformed function for a particular argument signature. `jax.pmap` maps a function across devices, but synchronization or gradient aggregation must be expressed with collectives such as `lax.psum` or `lax.pmean`. These transformations compose according to their types and semantics, and different orders can compute different quantities.

The same minimalist core delegates neural network abstractions to companion libraries (Flax, Haiku, Equinox) and optimization to Optax. This separation reflects the functional philosophy: the core provides transformations, while libraries build conventional abstractions on top. The trade-off is that production readiness depends not only on the transformation model, but also on the maturity of the surrounding libraries, export paths, and operational tooling for the target environment.

The functional constraints that JAX imposes become advantages in specific domains. Custom differentiation—higher-order gradients, custom vector-Jacobian product (VJP) and Jacobian-vector product (JVP) rules—composes cleanly because pure functions make differentiation rules predictable. Research on optimization algorithms benefits from transformations that let researchers manipulate

gradient computation as naturally as they manipulate data. Compilation-heavy accelerator workloads use XLA to extract more utilization when the program can be expressed in this functional style. Scientific computing with AD requirements benefits from functional purity that enables mathematical reasoning about code. JAX requires more upfront investment than PyTorch: the functional paradigm has a learning curve, state management requires explicit patterns, and debugging compiled code is harder than eager execution. Teams should choose JAX when its strengths align with project requirements, not as a default.

7.7.4 Framework trade-offs under measurement

The preceding sections described each framework’s design philosophy in qualitative terms: graph-first vs. eager-first, stateful vs. functional. The useful comparison is not which framework is fastest in the abstract, because that answer changes with model shape, batch size, hardware backend, and compiler configuration. The useful comparison is what each design lets the system see and optimize. Table 7.10 therefore maps TensorFlow, PyTorch, and JAX back to the three framework problems: execution visibility, differentiation model, and hardware abstraction path.

Table 7.10: Framework Design Trade-Offs: Each column reflects a distinct answer to the three core problems. TensorFlow makes deployment and graph capture central, PyTorch makes eager execution and debugging central, and JAX makes functional transformation and compiler visibility central. The table is a diagnostic map, not a benchmark ranking.

Aspect	TensorFlow	PyTorch	JAX
Graph Type	Static roots, dynamic front end in 2.x	Dynamic	Functional transformations
Programming Model	Imperative front end, graph capture path	Imperative	Functional
Core Data Structure	Tensor with framework-managed state	Tensor with framework-managed state	Immutable array
Execution Mode	Eager by default, graph for optimization	Eager by default	Trace and just-in-time compilation
Automatic Differentiation	Reverse mode over captured computation	Reverse mode over an eager tape	Forward and reverse transformations
Hardware Abstraction	Broad deployment runtimes and XLA paths	Native GPU path plus export/compile runtimes	XLA-centered accelerator compilation
Optimization Risk	Graph capture and operator coverage	Graph breaks after eager development	Purity, shape stability, and tracing

The measurement implication is straightforward: profile the constraint each framework makes most visible. In PyTorch, check whether eager dispatch or graph breaks dominate. In TensorFlow, check whether the captured graph covers the operators and deployment target. In JAX, check whether shapes and purity let XLA compile the program actually executed. A framework-level comparison or leaderboard can orient a decision, but it cannot replace profiling the specific workload on the target hardware.

The same simple network exposes how each design philosophy shapes the code. Listing 7.28 implements one neural network, a single linear layer mapping ten inputs to one output, across all three frameworks.

These three implementations solve the same mathematical problem but reveal distinct answers to the three problems. The differences are not cosmetic; they shape debugging workflows, deployment options, and optimization potential.

PyTorch binds state and computation together through class inheritance (`nn.Module`), solving the execution problem through eager evaluation: the graph builds as Python runs, making standard debuggers and control flow work naturally. Without graph capture, an optimizer cannot see the full computation before execution begins.

TensorFlow/Keras also executes eagerly by default in TensorFlow 2, while structured models can be captured with `tf.function` and exported to supported deployment runtimes. Portability still depends on operator coverage and conversion for each target.

JAX treats the model as a pure function²⁶ with immutable data and no internal state. In this

²⁶ | **Pure function:** Returns outputs determined by its inputs without relying on untracked side effects; JAX transformations are designed for pure functions. Ordinary Python effects may run during tracing but are not represented in the compiled computation, and impure code can fail or behave unexpectedly. Runtime printing or external effects require supported mechanisms such as `jax.debug.print` or callbacks; random draws require explicitly managed keys.

²⁷ | **Just-in-time (JIT) compilation:** Traces a function for an argument signature and lowers it to backend-specific executable code; the first call pays tracing and compilation costs; compatible later calls can reuse the cached executable. New shapes, dtypes, devices, or static arguments may trigger another compilation. Both compilation time and cached-call overhead depend on the program, backend, and hardware.

model, `grad`, `vmap` (automatic vectorization), and `jit` (just-in-time compilation²⁷) are composable transformations on stateless functions rather than methods attached to a mutable object system. The cost is explicit parameter management and a programming model that may be unfamiliar to engineers coming from stateful frameworks.

Listing 7.28: Framework Comparison: Hello World: The same simple neural network implemented in PyTorch (object-oriented), TensorFlow/Keras (declarative), and JAX (functional), illustrating each framework’s distinct design philosophy.

```
# PyTorch - Dynamic, Pythonic
import torch.nn as nn

class SimpleNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc = nn.Linear(10, 1)

    def forward(self, x):
        return self.fc(x)

# TensorFlow/Keras - High-level API
import tensorflow as tf

model = tf.keras.Sequential(
    [tf.keras.layers.Dense(1, input_shape=(10,))]
)

# JAX - Functional approach
import jax.numpy as jnp
from jax import random

def simple_net(params, x):
    return jnp.dot(x, params["w"]) + params["b"]

key = random.PRNGKey(0)
key_w, key_b = random.split(key)
params = {
    "w": random.normal(key_w, (10, 1)),
    "b": random.normal(key_b, (1,)),
}
```

No framework maximizes all three goals simultaneously, and the ten-line program makes the trade visible in source code rather than buried in compiler internals: where state lives, when the graph exists, and what the compiler is allowed to see. The underlying currency is compiler visibility vs. human iteration latency. Graph capture and ahead-of-time compilation can reduce runtime overhead and intermediate data movement; eager evaluation shortens the human iteration loop, outside the iron-law runtime equation, at the cost of compiler visibility until a graph-capture path is used; functional purity gives XLA more freedom to transform a stable traced program. Each philosophy also shapes code syntax, team workflows, debugging practices, and deployment pipelines, so migration costs grow with project-specific integrations and assumptions.

7.7.5 Runtime constraint map

The same principle extends beyond the three general-purpose frameworks. Runtime families differ because they remove different degrees of flexibility in exchange for compiler visibility, smaller binaries, or hardware-specific execution. Table 7.11 is therefore a constraint map: each row shows what the runtime gives up, what optimization path that enables, and which failure mode to test before committing.

Table 7.11: Runtime Constraint Map: Read each row left to right as one bargain, where a family gives up flexibility, gains an optimization path, and acquires a new way to fail. The middle column records where a family sits relative to an untuned eager baseline rather than a measured result, so latency, memory use, energy, and utilization still have to be measured for the specific model, hardware, precision, batch size, and compiler configuration.

Runtime family	Where it fits	Constraint anchor	Optimization path	Constraint to test
PyTorch eager	Research and iteration	Baseline latency; full Python/runtime footprint	Dynamic graphs, eager debugging	Dispatch overhead and missing graph view
Compiled PyTorch/TensorFlow	Server training and serving	Workload-dependent gains when graph capture is clean	Graph capture, fusion, layout planning	Operator coverage and graph breaks
TensorFlow Lite/Core ML	Mobile and edge inference	Target-specific mobile latency and package-size budgets	Quantization, static graphs, NPU delegates	Target-specific conversion constraints
TF Lite Micro/microTVM	Microcontroller inference	Target-specific microcontroller RAM and flash budgets	Static allocation, INT8 kernels	Selected operators and constrained memory
ONNX Runtime	Cross-framework serving	Backend-dependent; can match native runtimes for supported graphs	Standard graph format, execution providers	Export gaps and custom-operator fallbacks
TensorRT/TVM	Hardware-specialized inference	Target-dependent gains over untuned eager baselines	Kernel fusion, precision lowering, autotuning	Narrower target and conversion assumptions

The map reveals one systems pattern rather than a product ranking. Each move toward a narrower runtime trades flexibility for a more predictable execution plan. Specialized inference runtimes such as TensorRT and Apache TVM can deliver substantial latency gains when the model converts cleanly and the deployment target is known. Mobile and microcontroller runtimes reduce footprint by removing training machinery and relying on static graphs, quantization, platform delegates, or fixed memory arenas. A delegate is a runtime plugin that routes supported operators to a target accelerator and falls back when the operator is unsupported. The engineering question is always what was removed to make the optimization possible, because unsupported operators, dynamic shapes, or graph breaks can erase the expected advantage.

These efficiency gaps become hard constraints beyond the server room. A multi-fold latency gap between eager execution and a specialized inference engine is an optimization opportunity on a cloud GPU. On a microcontroller, a framework that exceeds the device’s memory capacity cannot run. The selection criterion shifts from raw latency to whether the framework fits inside the memory and runtime envelope.

7.8 Deployment Targets

As ML models move from cloud servers to edge devices, the efficiency gaps measured in section 7.7.4 transform from optimization opportunities into hard deployment constraints. Framework selection must reweight the three core problems dramatically at the edge. The execution problem shifts from choosing between eager and graph execution to fitting computation inside the target’s latency and memory budgets. The differentiation problem often disappears entirely, since edge devices run inference only. The abstraction problem intensifies as systems target ARM or x86 processors, mobile NPUs or edge TPUs, and microcontrollers with kilobytes of memory.

Table 7.12 continues the same constraint map across the cloud-to-edge spectrum: each row identifies the runtime assumptions that fit the target envelope.

Table 7.12: Deployment Target Constraint Map: Runtime assumptions, optimization levers, and binding constraints for each deployment tier, from cloud servers to microcontrollers.

Environment	Runtime assumptions that usually fit	Optimization lever	Binding constraint
Cloud/Server	Full training/serving frameworks	Graph compilation, batching, lower precision	Throughput, cost
Edge	Static graph or portable serving runtime	Static graphs, lower-precision kernels	Workload-specific latency and memory

Environment	Runtime assumptions that usually fit	Optimization lever	Binding constraint
Mobile	App-integrated runtime with delegates	Accelerator delegates, compact model formats	Battery, thermal, app size limits
Microcontroller (TinyML)	Tiny runtime with fixed allocation	Static allocation, small integer kernels	Tight RAM, no dynamic memory

Table 7.12 shows why deployment target is a hard constraint rather than a late packaging step. The Smart Doorbell KWS model from section 7.3.5 exemplifies the microcontroller tier: a runtime with a fixed memory arena and compact C/C++ footprint is not a preference but a condition for fitting the device. This constraint creates the practical framework problem that ONNX addresses: organizations often train in one environment but deploy into another whose runtime assumptions are stricter.

The ONNX²⁸ format addresses this fragmentation by enabling model portability across many runtimes (ONNX Contributors 2019): train in PyTorch, export through ONNX, and deploy through ONNX Runtime or a hardware-specific backend. TensorFlow Lite has its own conversion path rather than being a direct ONNX target in typical workflows. Standard interchange formats reduce manual conversion work when moving between development and production environments, but they do not eliminate compatibility testing, operator coverage gaps, or custom-kernel work. Figure 7.14 captures this hub-and-spoke interoperability model—notice how ONNX sits at the center, accepting models from the six tools shown on the left and making them available to ONNX Runtime and the compatible tools shown on the right. The compression and serving choices in chapter 10 and chapter 13 sit on top of this export boundary.

²⁸ | **ONNX**: The “fragmentation” ONNX addresses is that the framework used for model development may not match the runtime best suited to the deployment target. ONNX defines a hardware-agnostic graph representation that decouples the two, reducing the engineer-months of manual model conversion that would otherwise be required each time a deployment target changes. The accepted trade-off is that ONNX export can lose framework-specific optimizations or custom operators, requiring fallback implementations.




Figure 7.14: Framework Interoperability: The inward spokes represent exports into the shared ONNX format, while the outward spokes represent loads into compatible tools and runtimes; portability requires support on both sides of that boundary.

ONNX reduces the cost of framework fragmentation, but it does not eliminate the initial selection decision. The remaining question is how to choose a framework for a specific project’s constraints.

7.9 Framework Selection

Framework selection is a constrained optimization problem across the same three framework problems. The question is not which framework is “best”; it is which execution model, differentiation system, and abstraction path survive the project’s constraints.

7.9.1 The framework selection trade-off space

Framework selection involves three interconnected tensions. The first is between development velocity and production performance: eager execution prioritizes iteration speed, while graph compilation prioritizes runtime optimization. Research teams that need to test ten architecture variants per day cannot afford minutes of compilation between experiments; production teams that deploy a single model for months cannot afford the throughput penalty of eager dispatch. The optimal point shifts as a project moves through its lifecycle.

This velocity-performance tension leads directly to the second, flexibility versus optimization depth. Eager execution makes host-language control flow natural but limits compiler scope until a larger region is captured. Static graphs can represent data-dependent control flow through graph operations while exposing more of the program for fusion and hardware-specific code generation. As table 7.3 demonstrated, this trade-off cascades through memory management, utilization, and debugging workflows. It is not a single design decision but a system-wide constraint.

The flexibility-optimization tension, in turn, exposes a third: ecosystem breadth vs. specialization. General-purpose frameworks cover broad operation sets but often underperform specialized runtimes on workloads those runtimes can optimize deeply. TensorRT, TVM, and similar systems optimize for narrower deployment targets through fusion, precision selection, and hardware-specific scheduling (NVIDIA 2024c; Chen et al. 2018). ONNX bridges part of this gap through standardized interchange (ONNX Contributors 2019). Runtime specialization still has to be evaluated separately: the more a runtime specializes, the more it depends on conversion coverage, supported operators, and fallback behavior.

Systems Perspective 7.4: Framework selection constraints

Rather than seeking the “best” framework, effective selection first eliminates candidates that cannot satisfy hard constraints: deployment target, required operations, compiler path, and team expertise. Only then should soft preferences such as performance, development speed, and ecosystem rank the survivors.

The TensorFlow ecosystem illustrates how these axes interact concretely. Its three variants (TensorFlow, TensorFlow Lite, TensorFlow Lite Micro) trace a single design philosophy across progressively tighter constraints, a pattern that generalizes to any framework family. Table 7.13 traces the trade-offs.

Table 7.13: TensorFlow Variant Model Support: The four rows show what each variant can express, not how large it is. Training capability disappears across the tiers while inference and lower-precision tooling survive intact, which leaves operator coverage as the discriminating row: broad in TensorFlow, inference-focused in TensorFlow Lite, and reduced to the operators one model selects in TensorFlow Lite Micro. Binary size and memory footprint appear in table 7.15.

	TensorFlow	TensorFlow Lite	TensorFlow Lite for Microcontrollers
Training	Yes	Limited	No
Inference	Yes	Yes	Yes
Operator Coverage	Broad	Inference-focused	Model-selected subset
Native Lower-Precision Tooling	Yes	Yes	Yes

The principle is progressive constraint leading to progressive optimization: fewer supported operations enable smaller binaries, tighter memory budgets, and deployment-focused lower-precision execution. Three dimensions structure this analysis: model requirements define the supported operations, software dependencies define the runtime environment, and hardware constraints define the physical limits.

7.9.2 Framework selection criteria

Three dimensions structure systematic framework evaluation: what the model requires (supported operations and graph semantics), what the software environment provides (OS, memory management, accelerator delegation), and what the hardware physically permits (compute, memory, power).

Each dimension acts as a filter: hard constraints eliminate candidates, and soft preferences rank the survivors.

7.9.2.1 Model requirements

The first question is whether a framework can express the models a project requires. Examine table 7.13. Operator coverage narrows from full TensorFlow to TensorFlow Lite and then to the model-selected subset linked into TensorFlow Lite Micro. Each reduction narrows training capability and general-purpose operations while concentrating deployment tooling. The engineering principle is that algorithmic expressiveness and machine efficiency trade against each other. Fewer supported operations enable tighter code generation, smaller binaries, and hardware-specific optimization paths. This progressive constraint model applies to any framework family, not just TensorFlow. Layered on top of operator coverage is a separate axis with its own trade-off, whether the graph is captured statically before execution or assembled dynamically at runtime.

🔍 Systems Perspective 7.5: Dynamic vs. static computational graphs

The static-vs.-dynamic graph distinction (examined in section 7.3) has direct implications for model requirements analysis. Static graphs require control flow and dynamic behavior to be represented in forms understood by the graph runtime, which enables ahead-of-time optimization for deployment. Eager execution follows ordinary Python control flow directly, while graph-capture paths such as `torch.compile` and `tf.function` recover optimization opportunities when they can represent the executed program.

7.9.2.2 Software dependencies

Once model requirements are satisfied, the framework must integrate with the target software environment. Table 7.14 reveals how operating system requirements, memory management, and accelerator support vary across TensorFlow variants.

Table 7.14: TensorFlow Variant Capability Comparison: Capabilities of TensorFlow, TensorFlow Lite, and TensorFlow Lite Micro regarding operating system dependence, memory management, and hardware acceleration. Progressive constraint across variants enables selection by deployment context, from full-scale servers to resource-constrained edge devices.

	TensorFlow	TensorFlow Lite	TensorFlow Lite for Microcontrollers
Needs an OS	Yes	Yes	No
Model Access	Filesystem or service	Platform-specific storage mapping	Platform-specific compiled or mapped storage
Accelerator Support	Yes	Delegates	Optimized kernels and integrations

The key distinctions follow the same progressive constraint pattern. TensorFlow Lite Micro eliminates the OS requirement entirely, enabling bare-metal execution on microcontrollers (though it integrates with RTOSes like FreeRTOS and Zephyr when available). How each Lite runtime accesses a model depends on the platform and integration. TensorFlow Lite uses delegates for supported accelerators, while TensorFlow Lite Micro can use platform-optimized DSP or accelerator kernels without the standard delegate mechanism. Each software dependency removed is a deployment target gained.

7.9.2.3 Hardware constraints

Software compatibility alone does not guarantee deployment; the framework must fit within physical hardware limits. Table 7.15 sets out this final constraint dimension.

Binary size and memory footprint depend on the selected build, operators, delegates, model, and runtime configuration. Processor architecture support shifts from x86 processors, GPUs, and TPUs in data centers through Arm Cortex-A platforms at the mobile/edge tier to Arm Cortex-M processors, DSPs, and MCUs in embedded systems. These are not arbitrary engineering tiers—they mirror the

physical constraints (Light Barrier, power wall, memory wall) that carve the deployment spectrum into distinct paradigms (section 2.2). The engineering lesson generalizes beyond TensorFlow: every framework family that spans deployment tiers makes analogous trade-offs between capability and resource footprint, and the framework’s job is to make those trade-offs navigable rather than invisible.

Table 7.15: TensorFlow Hardware Targets: TensorFlow variants target progressively more constrained hardware, shifting from x86, GPUs, and TPUs to Arm Cortex-A platforms and then Arm Cortex-M processors, DSPs, and MCUs. Actual binary size and memory footprint depend on the build, selected operators, delegates, model, and runtime configuration.

	TensorFlow	TensorFlow Lite	TensorFlow Lite for Microcontrollers
Base Binary Size	Build- and package-dependent	Build- and delegate-dependent	Model- and operator-selection-dependent
Base Memory Footprint	Model-, runtime-, and allocator-dependent	Model-, delegate-, and allocator-dependent	Fixed arena plus runtime state
Optimized Architectures	x86, TPUs, GPUs	Arm Cortex-A, x86	Arm Cortex-M, DSPs, MCUs

7.9.2.4 Production-ready evaluation factors

Beyond this expressiveness-efficiency trade-off, technical specifications establish necessary but not sufficient conditions for selection. Production deployments also require evaluating migration cost, maintenance burden, and deployment reliability.

These hardware constraints cascade into tightly coupled performance trade-offs. Inference latency, memory footprint, power consumption, and hardware utilization depend on the model and target device rather than on the runtime family alone. Lower-precision execution can reduce memory, latency, and energy at the cost of numerical margin, and framework selection determines which optimization levers are available. Scalability introduces a further concern. Consistent deployment from microcontrollers to servers, smooth prototype-to-production transitions, and version management across deployed fleets all depend on the framework’s deployment toolchain. The three-dimension methodology illustrated here (model requirements, software dependencies, and hardware constraints) applies to any framework ecosystem, not just TensorFlow.

7.9.3 Development support and long-term viability assessment

Framework viability over a five-year production deployment depends on whether the ecosystem keeps the chosen execution, differentiation, and abstraction paths maintainable. Community composition matters because it determines which problems receive engineering attention: research-heavy ecosystems tend to improve experimentation and reproducibility first, production-heavy ecosystems tend to improve serving, monitoring, and compatibility first, and smaller specialized ecosystems tend to advance narrower mathematical or compiler capabilities faster than broad deployment tooling.

A framework’s practical utility often depends more on these surrounding paths than on the core tensor API. Model hubs, experiment trackers, serving runtimes, cloud ML services, and interchange formats can reduce lock-in or deepen optimization, but each also adds a dependency that must be maintained. These compounding effects make framework migration progressively harder: CI/CD pipelines, monitoring infrastructure, cloud integrations, and custom operators turn an API choice into an operational commitment. The measurable indicators of viability are therefore contributor diversity, backward compatibility track record, available hiring pool, and the cost of preserving an exit path through standardized formats such as ONNX, framework-agnostic data pipelines, and documented customizations.

The three core problems have so far appeared in isolation: execution, differentiation, and abstraction examined one at a time, with framework choices and selection criteria layered on top. A single training step is where the three problems collide. Tracing one end-to-end reveals how framework execution, differentiation, and hardware abstraction operate as one integrated system.

7.10 Anatomy of a Training Step

Eight executable Python statements make eager vs. graph execution, reverse-mode autodiff, tensor abstractions, and kernel dispatch interact inside a single training step. Tracing that step through

the PyTorch stack reveals how the execution, differentiation, and abstraction machinery operate simultaneously.

Listing 7.29 presents a minimal training iteration for a two-layer multilayer perceptron. Though only eight executable statements, this code exercises the entire framework stack: tensor allocation, kernel dispatch, autograd recording, gradient computation, and parameter updates. Tracing each phase reveals the three problems in action and connects the quantitative principles developed in section 7.1 to concrete execution.

Listing 7.29: Training Step Anatomy: A minimal training iteration for a two-layer MLP, exercising tensor allocation, kernel dispatch, autograd recording, gradient computation, and parameter updates.

```
# Single training step for a 2-layer MLP
x = torch.randn(32, 784, device="cuda") # Input batch
y = torch.randint(0, 10, (32,), device="cuda") # Labels

# Forward pass
optimizer.zero_grad()
h = torch.relu(x @ W1 + b1) # Hidden layer
logits = h @ W2 + b2 # Output layer
loss = F.cross_entropy(logits, y)

# Backward pass
loss.backward()

# Parameter update
optimizer.step()
```

Phase 1: Forward pass (solving the execution problem)

During the forward pass, when `h = torch.relu(x @ W1 + b1)` executes, PyTorch’s eager execution triggers immediate computation:

1. **Python dispatch:** The Python interpreter calls `torch.matmul`, which routes through PyTorch’s dispatcher to select the CUDA backend, adding microseconds-scale overhead before device work begins.
2. **Kernel selection:** cuBLAS selects an optimized GEMM kernel based on matrix dimensions ($32 \times 784 \times 256$). For these dimensions, it might choose a tiled algorithm optimized for L2 cache.
3. **Kernel launch:** The selected kernel is queued to the GPU’s command buffer, adding a few microseconds of launch overhead while the CPU continues immediately through asynchronous execution.
4. **GPU execution:** The kernel loads `W1` from HBM²⁹ to L2 cache, performs the matrix multiply in Tensor Cores when available, and writes the result back to HBM; for this small GEMM, the workload usually still lasts only microseconds.
5. **Autograd recording:** Simultaneously, PyTorch’s autograd engine records a `MmBackward` node on the tape, storing references to `x` and `W1` for gradient computation.

The bias addition and ReLU follow similar patterns, each adding a node to the autograd tape.

Phase 2: Backward pass (solving the differentiation problem)

Calling `loss.backward()` triggers a four-stage backward pass:

1. **Tape traversal:** The autograd engine traverses the recorded graph in reverse topological order.
2. **Gradient Computation:** For each node, it calls the registered backward function, where W_1 and W_2 are layer weight matrices that form part of the model parameters θ . Traversing in reverse, `CrossEntropyBackward` computes $\frac{\partial \mathcal{L}}{\partial \text{logits}}$ using the softmax derivative; `MmBackward` for W_2 computes $\frac{\partial \mathcal{L}}{\partial W_2} = h^T \cdot \frac{\partial \mathcal{L}}{\partial \text{logits}}$ along with $\frac{\partial \mathcal{L}}{\partial h}$; `ReluBackward` applies the ReLU derivative mask (zero where $h \leq 0$); and `MmBackward` for W_1 computes $\frac{\partial \mathcal{L}}{\partial W_1}$. The example does not request a gradient for the input tensor `x`.

²⁹ | **HBM (high bandwidth memory):** Provides 2–3 TB/s bandwidth on modern GPUs, making it the memory tier that most directly feeds accelerator arithmetic; HBM bandwidth determines whether operations are memory bound or compute bound, and its 80 GB capacity on an A100 sets the hard ceiling on total live training state. Weights, activations, gradients, optimizer state, and temporary workspace must fit during execution. When they do not, the framework must resort to offloading, selective recomputation, or placement across devices, each adding complexity to what the programmer perceives as a single `loss.backward()` call.

3. Gradient Accumulation: Gradients are accumulated into `.grad` attributes of leaf tensors.
4. Memory Management: Once a backward node's saved tensors are no longer needed and no other references retain them, the framework can release that state for memory reuse.

Together, these backward-pass stages turn the recorded forward graph into gradients while releasing intermediate state as soon as it is no longer needed.

Phase 3: Memory traffic analysis (the physics at work)

Applying equation 7.4 to this step, table 7.16 breaks down the FLOPs, memory traffic, and arithmetic intensity for each operation:

Table 7.16: Per-Operation Roofline Analysis: Analytical work and memory-traffic estimates for a two-layer MLP training step. The cross-entropy row assigns one operation-equivalent each to exponentiation, accumulation, and normalization per logit, while the backward row assumes twice the forward work and three times its memory traffic. These are modeling assumptions rather than measured instruction counts. MatMul operations have much higher arithmetic intensity than element-wise operations, while the simplified ReLU and cross-entropy models remain memory-sensitive.

Component	Modeled Work	Memory Traffic	Arithmetic Intensity
MatMul (x @ W1)	$2 \times 32 \times 784 \times 256 = 12.8 \text{ MFLOP}$	0.9 MB	13.7 FLOP/byte
ReLU	$32 \times 256 = 8.2 \text{ KFLOP}$	65.5 KB	0.125 FLOP/byte
MatMul (h @ W2)	$2 \times 32 \times 256 \times 10 = 163.8 \text{ KFLOP}$	44.3 KB	3.7 FLOP/byte
Cross-entropy (simplified)	$\sim 0.96 \text{ KFLOP}$	2.6 KB	0.4 FLOP/byte
Backward (assumed 2× forward)	$\sim 26 \text{ MFLOP}$	3.1 MB	8.3 FLOP/byte

The arithmetic-intensity column is the diagnostic column. Matrix multiplications reuse operands enough to move toward the compute roof, while ReLU and cross-entropy move too little work per byte to escape the memory and launch-overhead regime. This is why fusion and dispatch reduction matter even when the model is written in high-level tensor code.

The model estimates $\sim 39.1 \text{ MFLOP}$ of work and $\sim 4.2 \text{ MB}$ of memory traffic. To obtain ideal lower bounds, we divide these estimates by the A100 peak rates:

- $T_{\text{compute}} \approx 39.1 \text{ MFLOP} / 19.5 \text{ TFLOP/s FP32} \approx 2.0 \mu\text{s}$
- $T_{\text{memory}} \approx 4.2 \text{ MB} / 2.04 \text{ TB/s} \approx 2.1 \mu\text{s}$
- Assuming 12 ops kernel launches, $T_{\text{overhead}} \approx 12 \text{ ops} \times 15 \mu\text{s} \approx 180 \mu\text{s}$

Under the stated launch-count and per-launch assumptions, the modeled training step is overhead-bound. Small models are often sensitive to Python dispatch and kernel launch costs, which drives three common production practices that all serve to reduce the dispatch term rather than changing the model's mathematical work:

- `torch.compile` can speed favorable small-operation workloads by fusing operations and reducing kernel launches
- Batch size increases help amortize per-batch overhead
- Larger, compute-dense operations can amortize dispatch overhead

Phase 4: Hardware abstraction (solving the abstraction problem)

The same Python code runs on different hardware through abstraction layers, each pairing a backend library with a hardware-specific execution mechanism, as table 7.17 summarizes.

Each backend implements common tensor-operation semantics with hardware-specific optimizations. A single `loss.backward()` call can therefore trigger different code paths depending on the hardware. Floating-point precision, reduction order, kernel selection, and nondeterministic operations may produce numerical differences across backends, so equivalence must be evaluated within the framework's documented tolerances and determinism guarantees.

Table 7.17: Hardware Backends Behind a Single Code Path: Each backend implements the same tensor operations through a different library and execution mechanism, so identical Python maps onto distinct silicon.

Hardware	Backend library	Execution mechanism
CUDA GPU	cuBLAS (NVIDIA 2024a; Choquette et al. 2021)	GEMM kernels and CUDA streams for async execution
CPU	Intel oneMKL or OpenBLAS (Intel Corporation 2026a; OpenBLAS Project 2026)	Thread-level parallelism around optimized kernels
TPU	XLA (Google 2025)	Compilation to TPU-specific high-level optimizer (HLO) operations
Apple Silicon	Metal Performance Shaders	MPS backend

This detailed trace through a single training step demonstrates how deeply the three core problems interact. Even simple code exercises the full framework stack, and seemingly minor decisions—device placement, batch size, compilation mode—cascade through execution, differentiation, and abstraction layers in ways that are difficult to predict without systems-level understanding. The pitfalls that follow are the common failure modes when that systems view is missing.

Systems Perspective 7.6: The three problems in action

This trace reveals the three problems in concrete terms:

- Execution: Eager mode enables line-by-line debugging but incurs dispatch overhead
- Differentiation: Autograd tape records operations during forward, replays in reverse during backward
- Abstraction: Same code runs on GPU/CPU/TPU through backend-specific kernel implementations

Understanding this flow enables informed optimization: fuse operations to reduce overhead, use appropriate batch sizes, and match model scale to hardware capabilities.

7.11 Fallacies and Pitfalls

Framework selection involves subtle trade-offs where intuitions from conventional software engineering fail. The memory wall, kernel fusion constraints, and deployment target diversity create pitfalls that waste engineering effort and cause production systems to miss latency targets.

Fallacy: *“All frameworks provide equivalent performance for the same model architecture.”*

Engineers assume that ResNet-50 yields identical performance across frameworks since the mathematics is the same, forgetting that performance is an emergent property of algorithm-machine co-design. In production, implementation matters. Compiled execution can improve supported eager workloads, and hardware-specialized inference engines such as TensorRT or TVM can reduce latency relative to untuned eager baselines when conversion is clean. The difference arises from kernel fusion depth, graph optimization strategies, memory access patterns, precision support, and backend libraries that vary between frameworks. Organizations that assume equivalence can miss latency service level agreements and require costly last-minute framework migrations.

Pitfall: *Choosing frameworks based on popularity rather than project requirements.*

Engineers assume the most popular framework works for any project. In reality, deployment constraints dominate. A mobile runtime such as ExecuTorch or TensorFlow Lite and a microcontroller runtime such as TensorFlow Lite Micro make different assumptions about memory allocation, operator coverage, and hardware delegates; the former targets phones with GB-scale memory, while the latter often targets devices with hundreds of KB of RAM, commonly below 256 KB for small TinyML deployments. Teams that prototype edge applications without checking the final runtime can face memory bloat that exceeds device capacity or a late framework migration after development completes. Evaluate deployment targets per section 7.8 before selecting a training framework.

Fallacy: *“Framework abstractions eliminate the need for systems knowledge.”*

Engineers assume high-level APIs handle all optimization automatically. The Roofline Model (section D.2.1) proves otherwise. Element-wise operations such as ReLU perform little arithmetic per byte moved, so memory traffic and launch overhead can dominate even on accelerators with abundant compute. Section 7.3.1 explains how this imbalance can leave compute resources idle. Engineers must identify the limiting resource before choosing fusion, batching, layout, or precision optimizations.

Pitfall: *Ignoring vendor lock-in from framework-specific formats.*

Engineers assume framework migration is straightforward since models are “just math.” Converting TensorFlow SavedModel to PyTorch can require rewriting custom operations, validating numerical equivalence across large test suites, and retraining when operations lack exact equivalents. ONNX (section 7.8) improves portability, but custom operators, dynamic shapes, and backend-specific optimizations can still require manual work. Organizations that ignore this during initial framework selection face costly migrations when deployment requirements change or better frameworks emerge.

Fallacy: *Training framework choice is independent of production infrastructure.*

Engineers assume training framework choice is independent of deployment infrastructure. In practice, framework-infrastructure mismatches can impose substantial operational overhead. Some serving stacks provide atomic model swaps, while others require a process or container restart unless the deployment architecture adds version routing around them. Some frameworks and runtimes expose monitoring hooks directly; others require custom instrumentation. These are preview consequences of the serving and operations layers developed in chapter 13 and chapter 14; the local lesson is to evaluate the complete deployment stack during framework selection, including serving infrastructure, monitoring, and operational tooling.

Pitfall: *Increasing batch size without modeling activation memory.*

Engineers assume that if memory is available, larger batches always improve throughput. Larger mini-batches can amortize fixed dispatch overhead, but they also scale the activation memory that must remain live during training. A 7-billion-parameter model in FP16 consumes 14 GB, leaving 71.9 GB on an 80 GB A100 before accounting for other training state. Increasing batch size from 8 to 32 quadruples the batch-dependent activation footprint; transformer attention adds a large $\mathcal{O}(S^2)$ term in sequence length that makes each sample expensive. The resulting memory pressure can trigger recomputation strategies whose additional work reduces throughput despite the larger batch. Teams that blindly maximize batch size can achieve lower throughput than smaller batches that avoid these memory management pathways; section 8.3.2.1 formalizes the throughput target.

Fallacy: *Compilation overhead is negligible.*

Engineers assume compilation overhead is a one-time cost that pays off quickly. The hypothetical values in table 7.4 assign ResNet-50 higher compiled throughput and a compile window of 15 s to 30 s per graph change. Under those assumptions, a 10,000-image experiment with 10 code changes completes in 6.9 s in eager mode and 304.7 s in compiled mode, including recompilation overhead. The compiled workflow is therefore about 44.2× slower in this illustrative rapid-prototyping scenario. Teams that enable compilation during rapid prototyping can waste time waiting for recompilations that negate throughput gains.

Pitfall: *Using one execution policy for exploration and production.*

The right framework mode depends on the loop. During exploration, eager execution and small tests shorten feedback by avoiding repeated graph captures and recompilations. During production serving or long training runs, compilation can amortize setup across many requests or samples. The deciding factors are workload duration and stability: compilation pays when graph structure and shapes recur, whereas changing control flow or input shapes can invalidate cached artifacts and restart the amortization. Teams should choose the policy from measured iteration latency and steady-state throughput. One policy for both phases either slows research iteration or leaves production throughput unused.

7.12 Summary

Machine learning frameworks exist to solve three fundamental problems that would otherwise make deep learning impractical. The first is execution: deciding when and how computation runs. Frameworks navigate the trade-off between eager execution (immediate, debuggable, flexible) and graph execution (deferred, optimizable, deployable), while modern hybrid approaches like `torch.compile` attempt to provide both flexibility during development and optimization for production. The second is differentiation: computing gradients automatically. Frameworks implement reverse-mode automatic differentiation that applies the chain rule across supported operation compositions, producing derivatives subject to floating-point arithmetic and each operation's derivative rule. This software primitive makes training billion-parameter models possible through a `loss.backward()` call. The third is abstraction: targeting diverse hardware from a single interface. Frameworks provide tensor abstractions, intermediate representations, and runtime systems that hide hardware complexity while enabling efficient utilization across CPUs, GPUs, TPUs, and specialized accelerators.

These problems are interconnected and constrained by the iron law of performance (section 1.7): execution strategy determines dispatch overhead (L_{lat}), differentiation determines memory traffic (D_{vol}), and abstraction determines hardware utilization (η_{hw}). The memory wall can make data movement more expensive than computation, explaining why frameworks invest in kernel fusion, selective recomputation, lower-precision execution, and compilation pipelines.



Key Takeaways: The layer between math and hardware

- **Every framework solves three problems:** Execution determines how to run, differentiation how to train, and abstraction how to express. TensorFlow couples graph capture to broad deployment paths, PyTorch couples eager iteration to capture and export paths, and JAX organizes differentiation and compilation as composable function transformations. These are infrastructure commitments, not tooling preferences.
- **The memory wall drives optimization:** Compute capacity has grown far faster than memory bandwidth for decades, widening the cumulative gap that bounds data movement. Kernel fusion, selective recomputation, lower-precision execution, and data layout optimizations all target the data movement term (D_{vol}) in the iron law, not the compute term.
- **Compilation must amortize setup cost:** The compilation continuum principle in equation 7.2 quantifies when execution savings exceed compilation costs. Rapidly changing graphs often favor eager mode, while repeatedly executed stable graphs can benefit from progressive compilation from JIT to AOT. The dispatch overhead law in equation 7.4 explains why small operations can benefit disproportionately once compilation is reused.
- **Module abstractions automate state management:** Automatic parameter discovery, mode-dependent behavior, and hierarchical composition with serialization appear across major frameworks, enabling million-parameter optimization in a single optimizer step regardless of API syntax.
- **Deployment constraints should drive selection:** Inference engines, mobile runtimes, and microcontroller runtimes trade off operator coverage, memory allocation, compilation, and acceleration differently. Gaps between eager and specialized inference, or between server and microcontroller memory, are architectural constraints. Evaluate the deployment target before selecting a framework.


Understanding framework internals transforms how practitioners approach performance debugging and optimization. When a training job runs slower than expected, engineers who understand execution graphs can identify whether the bottleneck lies in eager-mode overhead, insufficient kernel fusion, or suboptimal memory layout. When deployment fails on target hardware, the compilation pipeline reveals whether the issue is operator support, quantization compatibility, or runtime configuration. This knowledge is essential for diagnosing and resolving performance issues in production systems.

A framework presents itself as a convenience, a cleaner way to write a model, and that is exactly what makes its influence easy to miss. Its real work is to translate mathematics into machine operations, and no translation is free: every choice it makes (eager or graph execution, when to fuse kernels, what precision to keep, how much to compile) shifts cost between the terms of the iron law rather than removing it. What looks like an API is therefore a standing decision about where the system will spend, made mostly before the engineer arrives. The framework cannot lighten the load the iron law names; it can only decide which of the three terms will carry it.

What's Next: From control room to power plant

Frameworks are the software substrate that translates abstract architectures into executable kernels. Computational graphs, autograd tapes, and kernel dispatch pipelines are the control room instruments: they give engineers visibility into and control over the training process. The machinery is built, but every budget decision it exposes remains unmade: which activations to checkpoint when memory runs short, how large a batch the hardware allows, and at what point a single device stops being enough. Chapter 8 is where those bills come due, scaling framework execution, differentiation, and memory-management machinery into the training systems that power modern AI.

Model Training




- 8.1 Training Systems Fundamentals
- 8.2 Iron Law of Training Performance
- 8.3 Mathematical Foundations
- 8.4 Pipeline Architecture
- 8.5 Identifying Bottlenecks
- 8.6 Pipeline Optimizations
- 8.7 Scaling Training Systems
- 8.8 Fallacies and Pitfalls
- 8.9 Summary

Purpose

Why can training a model cost orders of magnitude more than one inference?

Inference computes a forward pass as data flows through the network and a prediction emerges. Training adds a backward pass and an optimizer step, retains intermediate state, and repeats that work across many examples and optimization steps. Optimizers may also maintain momentum and variance estimates whose storage exceeds the model weights. The ratio depends on the model, workload, training duration, and inference volume, but one training run commonly costs orders of magnitude more than one inference. A lab that trains in three days can explore more designs than one that takes a month. A faster step is not faster training if it delays convergence, so throughput matters only relative to time to target quality. Across repeated experiments, even modest changes in convergence or utilization compound into substantial schedule and cost differences. The central systems challenge is to keep accelerators supplied with data and useful arithmetic without exhausting memory or losing time to synchronization. Which intervention helps depends on the binding compute, memory, data, or communication constraint, so profiling must precede scaling or optimization. Batch size, numerical precision, checkpointing, data loading, and parallelism must therefore be treated as coupled decisions rather than independent tuning knobs. For the systems engineer, training is where hardware and parallelism determine whether a run fits within memory, schedule, and budget, and whether the repeated optimization loop accelerates or stalls. In D·A·M terms, each precision and parallelism decision negotiates between the mathematics of optimization and the physics of execution.





Learning Objectives


- Explain training cost asymmetry using forward, backward, optimizer-state, data, and iteration costs
- Calculate FLOPs, activation memory, optimizer state, and dollar cost for neural network training
- Compare SGD, Adam, and AdamW by convergence behavior, memory overhead, and compute cost
- Diagnose compute-, memory-, and data-bound training bottlenecks with roofline analysis and profiling evidence
- Apply mixed precision, checkpointing, gradient accumulation, and FlashAttention to fit accelerator memory and throughput limits
- Design single-machine training pipelines with prefetching, overlap, batching, and systematic re-profiling
- Evaluate when to scale beyond one machine using memory, duration, communication, energy, and cost constraints

8.1 Training Systems Fundamentals

RUNNING a model once and training it from scratch live on opposite sides of the systems cost curve. Frameworks provide the execution substrate: computational graphs schedule operations, automatic differentiation computes gradients, and hardware abstractions target diverse accelerators. Those mechanisms make a single training step possible; training systems make it repeatable at scale. The same forward/backward/update loop may run many thousands or millions of times while retaining activations, feeding accelerators, and staying within practical memory, time, and budget limits.

This chapter uses an external order-of-magnitude estimate of approximately \$50,000 for GPT-2-scale training and a public estimate near \$100M for a GPT-4-class training run.¹ Neither figure was disclosed as an audited training bill. This **training cost asymmetry** reflects the difference between evaluating a trained model for a defined request and repeatedly executing forward, backward, and update computations across a training corpus.

For GPT-2, the chapter uses roughly 3.00×10^9 floating-point operations per token for a forward evaluation, following the parameter-based accounting associated with the architecture in the GPT-2 technical report (Radford et al. 2019). The total budget of 1.50×10^{21} FLOPs corresponds to roughly 167 billion token positions when forward and backward work is approximated as three times the forward cost. This accounting estimate is not a claim that GPT-2 executed that many complete inference requests. The scale makes training systems engineering a distinct discipline and shows how access to training infrastructure can constrain participation in AI development.



Training cost stays flat across scale, then explodes past the large-scale knee.

¹ | **Training cost scaling:** The roughly $2,000\times$ ratio between this chapter's two cost anchors illustrates how parameter counts, training-token budgets, hardware generations, and accelerator fleets can compound; it is not a controlled comparison of disclosed training bills. OpenAI's GPT-2 report documents the model and training setup but does not disclose training cost (Radford et al. 2019). Exact GPT-4 training details were also not disclosed, so the GPT-4-class figure is explicitly an industry estimate supported by public reporting and independent infrastructure analysis (Knight 2023; Patel and Wong 2023).



Definition 8.1: Training systems

Machine learning training systems are software-hardware systems that execute the iterative optimization loop (forward pass, loss computation, backward pass, and parameter update) to minimize a loss function over a training dataset.

1. **Significance:** Training memory cost is $8\times$ the inference memory cost per parameter in a standard mixed-precision Adaptive Moment Estimation (Adam) setup: a 7B-parameter model requires 14 GB (FP16 weights) + 14 GB (FP16 gradients) + 28 GB (FP32 master weights) + 56 GB (Adam first and second moments in FP32) = 112 GB at least, before accounting for activation storage. This multiplier is a major reason a model that runs inference on one GPU requires multiple GPUs for training.
2. **Distinction:** Unlike inference systems, which execute a single forward pass and discard intermediate activations, standard backpropagation stores the tensors needed by the backward pass, creating a footprint that grows with model depth and batch size.

3. **Common pitfall:** A frequent misconception is that training failures are compute problems. A common failure is an out-of-memory (OOM) error, a memory-management problem: tensors saved for backpropagation accumulate during the forward pass and can exhaust memory before the first gradient is computed.

Three characteristics distinguish training workloads from general-purpose computing:

- **Computational intensity:** The 1.50×10^{21} FLOPs budget spread over days of wall-clock time demands sustained PFLOP/s-scale throughput from hardware whose realized large-model throughput is often far below theoretical peak (Narayanan et al. 2021; Chowdhery et al. 2022).
- **Memory pressure:** Storing 1.5B weights requires 6 GB in FP32; the Adam optimizer adds two state tensors per parameter, consuming another 12 GB. Activation memory adds a workload-dependent footprint governed by architecture, sequence length, batch size, precision, and checkpointing, and the combined state can exceed a single accelerator’s capacity.
- **Data dependencies:** Each gradient update depends on the result of the previous one, creating sequential bottlenecks that limit how much parallelism the system can exploit.

Each challenge points to a different kind of fix. Computational intensity pushes the system toward higher accelerator utilization and lower-precision arithmetic. Memory pressure calls for techniques such as gradient checkpointing,² a specific application of *rematerialization* (discarding and recomputing intermediate values to save memory, from chapter 7) that trades recomputation for reduced activation storage, and mixed-precision training,³ which reduces the memory footprint of weights and activations. Data dependencies motivate pipeline designs that overlap computation with data movement, building directly on the data loading throughput optimized in chapter 4 so the accelerator never sits idle waiting for the next batch. The current chapter focuses on single-machine and single-node multi-GPU training; scaling to hundreds of machines across network boundaries introduces communication and fault tolerance challenges beyond this chapter’s scope.

8.1.1 The staged system pipeline: Identifying “accelerator bubbles”

A training system is not a single loop; it is a **staged system pipeline**. Reaching high accelerator utilization requires analyzing training as a factory floor where four distinct stages coordinate to keep the ALUs busy:

1. **Data Loading & Preprocessing** (CPU/Storage): Fetching raw bits from Non-Volatile Memory Express (NVMe), decoding, and augmenting on CPU cores.
2. **Host-to-Device Transfer** (PCIe): Moving the processed batch over the PCIe bus into GPU memory via direct memory access.
3. **Forward/Backward Pass** (Accelerator): Propagating activations and gradients through layers on the GPU.
4. **Parameter Synchronization** (Interconnect): Exchanging gradients between accelerators over the available fabric. NVLink provides up to 900 GB/s in the H100 configuration used later, while other systems may use PCIe or a network interconnect.

Any mismatch in the throughput of these stages creates **accelerator bubbles**, intervals in which some accelerator resources are idle or underutilized while waiting for another stage. A systems engineer reduces these bubbles through techniques such as asynchronous prefetching and pipeline overlap, aiming to make the next batch available before the current step needs it. Section 8.5 develops the quantitative tools for measuring and reducing these bubbles.

The chapter follows that dependency chain. The iron law of training performance comes first as a specialized application of the general iron law (section 1.7), separating total operations, peak throughput, and utilization. That equation provides the accounting system for the mathematical foundations that follow: neural-network computation as a workload, optimizer behavior, backpropagation mechanics, and arithmetic intensity. Once the costs are visible, the chapter turns to the training pipeline itself, where data loading, forward pass, backward pass, and parameter updates each constrain the next. The optimization sections then target the terms exposed by the accounting: mixed-precision training, FlashAttention, gradient accumulation, checkpointing, and

² | **Gradient checkpointing (activation checkpointing):** Backpropagation retains the forward values required by each backward rule. Checkpointing saves selected values and recomputes others during backward, trading additional work for lower activation memory. The $\sqrt{N_L}$ checkpoint placement and its associated recomputation cost describe the classical schedule analyzed by T. Chen et al. (2016) rather than every checkpointing implementation. Section 8.6.5.2 works through that schedule for GPT-2’s layer count.

³ | **Mixed-precision training:** Uses lower precision for selected computation and storage while retaining higher precision where accumulation or numerical range requires it. A common FP16 recipe maintains FP32 master weights and uses loss scaling to reduce gradient underflow (Micikevicius et al. 2017). BF16 (“Brain Floating Point,” from Google Brain (Wang and Kanwar 2019)) matches FP32’s 8-bit exponent range and therefore usually avoids dynamic loss scaling, though the exact state and accumulation precision depend on the optimizer and implementation.

data prefetching. Scaling beyond one accelerator comes last, after the single-machine levers expose communication overhead as the next bottleneck.

Before formalizing the iron law, consider how these constraints interact in practice. The theoretical framework matters because failures are expensive: a single gradient explosion can erase days of computation worth thousands of dollars.



War Story 8.1: The PaLM loss spikes (2022)

Context: PaLM’s 540-billion-parameter run used 6,144 Tensor Processing Unit (TPU) v4 chips across two 3,072-chip Pods connected through Google’s data-center network and suffered severe training-instability spikes absent from smaller runs (Chowdhery et al. 2022).

Failure mode: The training loss spiked roughly 20 times despite gradient clipping. The spikes occurred at highly irregular intervals, sometimes late in training, never appeared in smaller models, and had no identified principled mitigation.

Impact: Each recovery discarded recent progress on an expensive 6,144-chip run, so reproducible checkpoints and repeatable data order were prerequisites for replay.

Response: Engineers restarted training from an uncorrupted checkpoint roughly 100 steps prior to each spike and skipped approximately 200–500 data batches around the failure window.

Systems lesson: Operational fault tolerance directly governs the utilization factor η_{hw} in the iron law of training performance (equation 8.1). Automated loss-spike detection, checkpoint cadence, and batch skipping are core system mechanisms required to protect cluster compute efficiency.

8.2 Iron Law of Training Performance

PaLM’s loss spikes illustrate a broader point: instability at frontier scale is not a rare accident but a predictable consequence of pushing computational intensity, memory pressure, and data dependencies simultaneously. Each of those three characteristics is individually manageable, yet their interactions produce failure modes that operational recovery alone cannot prevent. Diagnosing which factor dominates a given failure, and quantifying the cost of each, requires a formal decomposition of training time into its physical constituents.

The iron law provides exactly that organizing framework: it decomposes training time so that every optimization technique maps to a specific term in the equation. This is a specialized application of the general iron law of ML systems introduced in section 1.7, focused specifically on maximizing computational throughput.



Definition 8.2: The iron law of training performance

The iron law of training performance expresses training time through the workload’s operation count and its effective achieved throughput:

$$T_{\text{train}} = \frac{O}{R_{\text{peak}} \times \eta_{hw}} \quad (8.1)$$

Here η_{hw} is an effective end-to-end utilization factor, the achieved model-operation rate divided by the relevant hardware peak. Data stalls, communication, synchronization, launch overhead, and inefficient kernels all reduce this factor. Prefetching and overlap can hide part of those costs but do not guarantee their removal.

1. Significance: The three factors identify three distinct optimization levers: O (reducible by algorithmic changes, fewer training tokens, or later model-compression methods such as pruning and distillation), R_{peak} (changed by hardware and the selected precision), and η_{hw} (the effective utilization factor and primary systems target; Narayanan et al. (2021) measured 44 to 52 percent of theoretical peak throughput across GPT models trained on

A100 clusters). Chapter 10 defines the compression techniques; here they serve only to show where such methods enter the accounting.

2. **Distinction:** The general iron law exposes data movement, computation, and latency separately. This training form folds every source of lost throughput into η_{hw} , making it useful for budget accounting but insufficient for diagnosing which underlying term is responsible.
3. **Common pitfall:** A frequent misconception is that η_{hw} is fixed by hardware. System efficiency is a pipeline property: memory bandwidth saturation, kernel launch overhead, and synchronization barriers each reduce η_{hw} independently, and diagnosing which factor dominates requires profiling rather than reading hardware specs.

Equation 8.1 reveals three levers for improvement: reduce total operations through algorithmic innovation, increase peak throughput through hardware utilization, or improve utilization through better pipeline orchestration. Each optimization technique in this chapter pulls one or more of these levers, as summarized in table 8.1.

Table 8.1: Iron law optimization mapping: Optimization techniques mapped to iron law terms. Understanding which term a technique affects guides optimization strategy selection.

Technique	Term Affected	Mechanism
Mixed Precision (FP16/BF16)	Peak Throughput $\uparrow$	Uses supported lower-precision hardware paths
Data Prefetching	Utilization $\uparrow$	Reduces accelerator idle time waiting for data
Gradient Checkpointing	Memory $\downarrow$, Operations $\uparrow$	Reduces saved activations by recomputing selected values
Gradient Accumulation	Memory $\downarrow$, Synchronization $\downarrow$	Builds an effective batch from serialized micro-batches
Operator Fusion	Memory Traffic $\downarrow$, Overhead $\downarrow$	Avoids intermediate transfers and reduces kernel launches
FlashAttention	Memory Traffic $\downarrow$, Utilization $\uparrow$	Same asymptotic FLOPs, much lower HBM IO; backward recomputation may add FLOPs

A caveat: the iron law focuses on execution efficiency—how fast the hardware processes a given workload. It does not capture data-side factors such as data quality, dataset size, or curriculum design, which affect how many total operations O are needed to reach a target accuracy. A cleaner dataset or a better data mix can reduce the number of epochs required, shrinking O without touching hardware at all. Holding the workload fixed, the question is how to execute it as fast as possible.

Actual training throughput remains below the arithmetic peak whenever kernels, data delivery, or communication leave resources idle. Scaling to multiple accelerators introduces additional communication overhead, a trade-off quantified in section 8.7.

Checkpoint 8.1: The physics of training

Training speed is governed by the utilization of hardware peaks.

Utilization gap

- ☐ **Peak vs. real:** why does sustained end-to-end utilization remain below the arithmetic peak?
- ☐ **Utilization losses:** which pipeline stalls introduced earlier reduce η_{hw} , and which can prefetching or overlap hide?
- ☐ **Order of magnitude:** a frontier run needs on the order of 10^{23} FLOPs at a utilization η_{hw} well below 1. Estimate which lever shortens wall-clock time more, doubling R_{peak} or doubling η_{hw} , and why the iron law says the two are not interchangeable in practice.

The iron law provides a static framework for reasoning about training performance, but the history of deep learning reveals how the binding constraint has shifted over time as hardware and algorithms co-evolved. In 1986, Rumelhart and colleagues popularized backpropagation for multilayer neural

networks (Rumelhart et al. 1986). In 2012, AlexNet trained an ImageNet classifier in five to six days on two GPUs (Krizhevsky et al. 2012), demonstrating how well neural-network parallelism mapped to graphics hardware. By 2017, transformers (Vaswani et al. 2017) shifted attention toward large-scale sequence modeling and high-throughput accelerator kernels. GPT-3 in 2020 consumed about 3.14×10^{23} FLOPs (Brown et al. 2020), making utilization (η_{hw}) critical. By 2023, training efficiency improved through the techniques examined in this chapter: FlashAttention reduces memory traffic while improving η_{hw} ; gradient checkpointing trades additional O for memory capacity; mixed precision increases R_{peak} . Each innovation addresses a specific iron-law bottleneck.

8.2.1 Running example: Training GPT-2

In inference, GPT-2 serves as the Bandwidth Hog lighthouse. Here it recurs as the chapter’s stable workload for testing each training cost and optimization.



Lighthouse 8.1: Lighthouse example: Training GPT-2

Context: GPT-2 (1.5 billion parameters) serves as the primary case study because it is large enough to expose meaningful compute, activation, and optimizer-state constraints while remaining easier to analyze than frontier-scale models. Whether it requires distributed training depends on precision, optimizer state, activation memory, hardware capacity, and the target training time. Table 8.2 maps each model property to the systems constraint it creates:

Table 8.2: GPT-2 (1.5 Billion Parameters) Lighthouse Model Specifications: Parameter count, architecture depth, dataset size, and total training compute mapped to their systems implications. Each row pairs a model property with the engineering constraint it creates—memory footprint for weights, activation pressure from pipeline depth, I/O throughput requirements, and parallelization demand.

Property	Specification	Systems Implication
Parameters	1.5B (XL)	Requires ~3 GB (FP16) or ~6 GB (FP32) for weights alone.
Architecture	48 Layers, 1600 Dim	Deep pipeline creates heavy activation memory pressure.
Dataset	WebText (40 GB)	Preprocessing, shuffling, and delivery must sustain the selected training configuration.
Compute	$\sim 1.50 \times 10^{21}$ FLOPs	Training duration depends on achieved throughput and accelerator count.

Mechanism: Training GPT-2 couples a total operation count of $O \approx 6 \cdot P \cdot D_{tokens}$, or 1.50×10^{21} FLOPs, with a strict memory footprint $M_{state} = M_{weights} + M_{grads} + M_{opt} + M_{act}$. Balancing execution time $T_{train} = \frac{O}{R_{peak} \cdot \eta_{hw}}$ requires orchestrating data ingestion, activation memory, and compute throughput within accelerator high-bandwidth memory (HBM) capacity.

Not all training workloads are compute bound. Recommendation models like DLRM are dominated by massive embedding tables (100B to 10T parameters, mostly embeddings) that make them memory bandwidth bound rather than compute bound. For such workloads, the first scaling problem is often capacity: splitting the embedding tables across devices so the model fits at all. The remainder of this chapter focuses on dense, compute-intensive training using GPT-2 as the primary worked example.

Training systems occupy a critical position in the machine learning pipeline: they consume prepared data from upstream engineering (chapter 4) and produce trained model artifacts that later systems must deploy and monitor. Data quality directly impacts training stability, while training efficiency determines iteration velocity during model development. The same pressure appears at three scales. At data scale, petabyte datasets require efficient I/O pipelines and distributed storage. At model scale, billion-parameter models force the system to decide whether to replicate the model across batches with data parallelism⁴ or split the model across devices with model parallelism.⁵ At infrastructure scale, coordinating thousands of accelerators introduces communication overhead that can dominate training time. These challenges are why the workflow contracts from chapter 3 matter during training: orchestration decisions shape both scientific iteration and systems cost.

⁴ | **Data parallelism:** Replicates the full model on every device, splitting only the data. Each device computes gradients independently, then an AllReduce operation synchronizes them—adding communication volume proportional to model size at every step. This synchronization tax limits scaling efficiency: doubling accelerators rarely halves training time once communication becomes the bottleneck.

⁵ | **Model parallelism:** Partitions the model’s layers across devices when it exceeds single-device memory. Activations must transfer between devices at every partition boundary, and naive partitioning creates “pipeline bubbles” where downstream devices idle while waiting. Microbatch pipelining, as in GPipe and PipeDream, recovers much of this lost efficiency (Huang et al. 2019; Narayanan et al. 2019).

GPT-2’s 40 GB WebText corpus sits at the lower end of this data scale spectrum. Larger corpora can change the data path qualitatively.

Systems Perspective 8.1: The 10 GB to 10 TB scale factor

In a memory-resident regime, a corpus may fit in system RAM and avoid repeated storage reads after caching, though preprocessing and transfer costs remain. In a streaming regime, the corpus exceeds local memory and the system must orchestrate storage reads, preprocessing, caching, and delivery throughout training. The appropriate design may include multiple workers, prefetching, local caches, or distributed storage, depending on which stage limits throughput.

Scale therefore changes more than the amount of data; it transforms the system’s physics.

These scaling challenges translate into concrete workflow requirements. Training workflows consist of interdependent stages—data preprocessing, forward and backward passes, and parameter updates—extending the neural network concepts from chapter 5. System constraints often dictate performance limits: accelerators with high compute-to-bandwidth ratios are frequently bottlenecked by memory bandwidth, where data movement between memory hierarchies is slower than the computations themselves (Hennessy and Patterson 2017). In distributed setups, synchronization across devices introduces additional latency, with interconnect performance (NVLink, InfiniBand) critically affecting throughput.⁶

The same hardware-software boundary that frameworks exposed in chapter 7 is central here. Mixed-precision training emerged from recognizing that Tensor Core hardware could accelerate reduced-precision arithmetic. Gradient checkpointing arose from memory capacity constraints. Training systems engineering is the work of matching those algorithmic choices to the physical limits of the machine.

These scaling challenges share a common thread: every bottleneck traces back to the cost of specific mathematical operations. Dense matrix multiplication is a well-studied systems kernel (Goto and Geijn 2008), while activation functions consume memory bandwidth and optimizer states increase the resident memory footprint. Designing systems to execute these operations at scale begins with accounting for each cost separately.

8.3 Mathematical Foundations

Matrix multiplication is just $\mathbf{C} = \mathbf{AB}$ in notation, but training GPT-2 repeatedly executes this operation on matrices too large to remain entirely in the fastest memory. Tensor Cores accelerate eligible matrix operations, while choices such as rectified linear unit (ReLU) versus sigmoid affect the cost, fusion opportunities, and memory behavior of the surrounding elementwise kernels. Chapter 5 established *what* neural network operations compute and *why* they enable learning; the systems question is *what* they cost in FLOPs, memory, and bandwidth when those operations execute at scale.

Four dimensions structure this cost analysis:

- FLOP counts of the matrix operations that dominate dense neural-network training.
- Memory requirements for storing activations and optimizer states simultaneously.
- Bandwidth demands that determine whether operations are compute bound or memory bound.
- Arithmetic intensity classifications that guide optimization strategy selection.

Together, these dimensions provide the vocabulary for analyzing the computational intensity, memory pressure, and data dependencies introduced in section 8.1.

8.3.1 Neural network computation

Neural network training consists of repeated matrix operations and nonlinear transformations. These operations are conceptually simple but create the system-level challenges that dominate modern training infrastructure. The introduction of backpropagation⁷ by Rumelhart et al. (1986) and the

⁶ | **Transformer training interconnect sensitivity:** Self-attention’s $\mathcal{O}(S^2)$

memory and compute scaling increases activation size with sequence length. In pipeline parallelism, activations cross devices at model-partition boundaries, while data parallelism synchronizes gradients whose volume scales with parameter count. The resulting communication cost depends on the parallelism strategy, topology, message size, and degree of overlap, making interconnect choice a first-order training-system decision.

⁷ | **Backpropagation provenance:** The algorithm was independently derived by Lin-nainmaa (1970) for automatic differentiation of computer programs and by Werbos (1974) in a Harvard PhD thesis on economic modeling—over a decade before Rumelhart et al. (1986) popularized it for neural networks. This delay between derivation and broad adoption recurs in ML systems history: attention mechanisms predated transformers by decades, but the 2017 Transformer showed that attention-heavy models could train efficiently on contemporary GPU systems; later TPU and GPU clusters enabled much larger deployments. A framework’s backward pass traverses the computational graph in reverse topological order and applies the chain rule; independent subgraphs may expose parallel work during that traversal.

8 | BLAS (basic linear algebra subprograms): The original BLAS specification standardized reusable Fortran-callable vector operations (Lawson et al. 1979); later BLAS extensions added matrix-vector and matrix-matrix routines, giving the familiar Level 1, Level 2, and Level 3 hierarchy (Dongarra et al. 1988). Training is dominated by Level 3 operations precisely because their high arithmetic intensity— $\mathcal{O}(n)$ FLOP/byte—can sustain high compute utilization. Frameworks commonly dispatch dense matrix multiplication to optimized libraries such as cuBLAS and oneDNN or to compiler-generated kernels.

9 | Strassen’s algorithm: Achieves $\mathcal{O}(n^{2.807})$ by replacing one of eight sub-multiplications with additions, but training systems rarely benefit. The recursion disrupts the regular memory-access patterns that Tensor Cores exploit, and accumulated rounding errors across billions of training iterations can destabilize convergence. In practice, mainstream general matrix multiplication libraries such as cuBLAS leave the dominant training matrix multiplications to highly optimized blocked $\mathcal{O}(n^3)$ kernels with superior hardware utilization rather than recursive Strassen variants.

development of efficient matrix computation libraries such as Basic Linear Algebra Subprograms (BLAS)⁸ (Dongarra et al. 1988) laid the groundwork for modern training architectures.

8.3.1.1 Mathematical operations in neural networks

Forward propagation, in its simplest case, involves two operations: matrix multiplication and activation function application. Matrix multiplication implements the linear transformation at each layer. At layer ℓ , the computation can be described as (following the row-vector convention established in chapter 5):

$$\mathbf{A}^{(\ell)} = f(\mathbf{A}^{(\ell-1)}\mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)})$$

where:

- $\mathbf{A}^{(\ell-1)}$ represents the activations from the previous layer (or the input layer for the first layer), with each row being a sample in the batch,
- $\mathbf{W}^{(\ell)} \in \mathbb{R}^{n_{\ell-1} \times n_{\ell}}$ is the weight matrix at layer ℓ , which contains the parameters learned by the network,
- $\mathbf{b}^{(\ell)}$ is the bias vector for layer ℓ ,
- $f(\cdot)$ is the activation function applied elementwise (for example, ReLU, sigmoid) to introduce nonlinearity.

8.3.1.2 Matrix operations

Section 5.4.2.2 established that forward propagation reduces to chains of matrix multiplications, and section 6.9.1 catalogued the computational primitives—general matrix multiplication, convolution, and dynamic attention—that every architecture shares. Training amplifies these patterns: each operation executes not once but billions of times, and each forward pass is paired with a backward pass that roughly doubles the computational cost. Understanding which matrix operations dominate—and how their shapes change between forward and backward passes—reveals *why* specific system designs and optimizations emerged for training.

Matrix multiplication dominance has driven both algorithmic and hardware innovations. Early neural network implementations relied on standard CPU-based linear algebra libraries, but the scale of modern training demanded specialized optimizations. Strassen’s algorithm⁹ reduced the naive $\mathcal{O}(n^3)$ complexity to approximately $\mathcal{O}(n^{2.807})$ (Strassen 1969), and contemporary hardware-accelerated libraries like cuBLAS (NVIDIA 2024a) continue pushing computational efficiency limits.

This computational dominance has driven system-level optimizations: blocked matrix computations that parallelize across multiple units, and memory hierarchies designed for the access patterns of both forward and backward passes. As neural architectures grew, weight and activation matrices both had to remain accessible for backpropagation, and hardware evolved to serve these dense multiplication patterns within growing memory budgets. To illustrate the scale of these operations concretely, consider the attention layer computations in our GPT-2 lighthouse model.

A single GPT-2 layer makes the scale of these computations concrete.



Napkin Math 8.1: GPT-2 attention layer computation

Each GPT-2 layer performs attention computations that exemplify dense matrix multiplication demands. For one GPT-2 transformer layer (all heads combined) with batch size $B = 32$, sequence length $S = 1024$, and hidden dimension $d = 1600$:

Query, Key, Value Projections (the three linear transformations that create attention inputs—3 separate matrix multiplications):

$$\begin{aligned} \text{FLOPs} &= 2 \times 3 \times (B \times S \times d \times d) \\ &= 2 \times 3 \times (32 \times 1024 \times 1600 \times 1600) \approx 503 \text{ billion FLOPs} \end{aligned}$$

Attention matmuls ($\mathbf{QK}^T$ and $\mathbf{A}_{\text{attn}}\mathbf{V}$):

$$\begin{aligned} \text{FLOPs}_{\mathbf{QK}} &= 2 \times B \times N_{\text{heads}} \times S \times S \times d_{\text{head}} = 107.4 \text{ billion FLOPs} \\ \text{FLOPs}_{\mathbf{A}_{\text{attn}} \times \mathbf{V}} &= \text{FLOPs}_{\mathbf{QK}}; \quad \text{pair total} \approx 214.8 \text{ billion FLOPs} \end{aligned}$$

Output projection (attention output $\rightarrow$ hidden):

$$\text{FLOPs} = 2 \times B \times S \times d \times d \approx 168 \text{ billion FLOPs}$$

Feed-Forward Network (Two linear transformations with expansion factor 4):

$$\text{FLOPs} \approx 16 \times B \times S \times d^2$$

Computation Scale

- Per-layer forward total (QKV 503 + attention 214.8 + output proj 168 + FFN): ~2228.01 GFLOP
- With 48 layers in GPT-2: ~320.8 TFLOP per training step
- Across this illustrative 50,000 steps scenario: ~16041.7 PFLOP

Systems insight: A V100 GPU (125 TFLOP/s peak with Tensor Cores, 15.7 TFLOP/s without) would require 2.6 s for the modeled attention-plus-FFN training-step estimate at 100 percent utilization (theoretical peak; practical throughput would be lower). Reaching a 180 to 220 ms training step is therefore already a multi-accelerator lower-bound problem: ideal arithmetic alone requires roughly 12 to 15 V100s, and practical systems need additional headroom for utilization losses, communication, and pipeline overhead.

These FLOP counts are not academic bookkeeping. They are the compute term of the iron law made concrete, and they explain why training cost scales as a predictable function of model architecture and sequence length rather than as an unpredictable emergent property.

8.3.1.3 Matrix-vector and batched operations

Not all operations in neural networks involve large matrix-matrix multiplications. Normalization layers, bias additions, and certain recurrent computations involve matrix-vector operations instead. Although computationally simpler than matrix-matrix multiplication, these operations present distinct system challenges: they exhibit lower hardware utilization due to their limited parallelization potential. A single vector provides insufficient work to keep thousands of accelerator cores busy simultaneously. This characteristic influences both hardware design and model architecture decisions, particularly in networks processing sequential inputs or computing layer statistics.

Recognizing the limitations of matrix-vector operations, the introduction of batching transformed matrix computation in neural networks. By processing multiple inputs simultaneously, training systems convert matrix-vector operations into more efficient matrix-matrix operations. This approach improves hardware utilization but increases memory demands for storing intermediate results. Modern implementations must balance batch sizes against available memory, leading to specific optimizations in memory management and computation scheduling.

The progression from matrix-vector to batched matrix-matrix operations helps explain the hardware design choices in modern accelerators. Google's first TPU, designed for inference, incorporated a specialized matrix unit and memory hierarchy for dense operations (Jouppi et al. 2017). At a different scale, GPT-3 training used distributed accelerators to execute the matrix-matrix multiplication patterns exposed by large batches (Brown et al. 2020).

Systems Perspective 8.2: Why GPUs dominate training

The matrix operations described in section 8.3.1.2 directly explain data-center training hardware architecture. GPUs became central to large-scale training for three reasons:

- Matrix multiplication's independent element calculations map well to thousands of GPU cores (NVIDIA A100 has 6,912 CUDA cores).
- Specialized hardware units like Tensor Cores provide substantially higher peak throughput for supported lower-precision matrix operations.

- Blocked matrix computation patterns enable efficient use of GPU memory hierarchy (L1/L2 cache, shared memory, global memory).

When a later illustrative GPT-2 scenario assumes a $2.4\times$ V100 mixed-precision speedup, Tensor Core acceleration of eligible matrix multiplications is one source of the gain. The realized end-to-end ratio also depends on nonmatrix operations, memory traffic, and input and communication overhead.

Matrix multiplications dominate training compute, but neural networks require more than linear transformations. Between each layer's matrix operations, activation functions introduce the nonlinearity that enables networks to learn complex patterns. These functions appear computationally trivial compared to matrix multiplication, yet their implementation characteristics affect training efficiency in ways that matter at scale.

8.3.1.4 Activation functions

Activation functions like sigmoid, tanh, ReLU, and softmax introduce nonlinearity, but their implementation characteristics also shape training system performance. From a systems perspective, the choice of activation function determines computational cost, hardware utilization, and memory access patterns during backpropagation.

The critical question for ML systems engineers is not *what* these functions do mathematically, but *how* their cost behaves at scale. The benchmarks and trade-offs that follow build toward one systems thesis. Because elementwise activations move many more bytes than they compute, their kernels are often bounded by memory bandwidth rather than arithmetic, so the per-operation differences below matter less than their magnitudes first suggest.

Activation costs can accumulate across many elements and layers, but their contribution to end-to-end time depends on tensor shape, fusion, implementation, precision, and hardware. Figure 8.1 preserves one illustrative Apple M2 timing snapshot from the chapter's development history. Its benchmark configuration was not recorded, so the values should not be interpreted as reproducible evidence or transferred to other systems. On accelerators, activation kernels are often limited by memory traffic, fusion boundaries, and special-function throughput rather than by scalar CPU instruction latency. The same activation may run as a standalone kernel, fuse with a neighboring operator, or use a compiler-selected approximation. The plotted scalar timing therefore cannot predict its cost inside a compiled training step.




Figure 8.1: Illustrative Activation Timing Snapshot: A legacy Apple M2 timing snapshot compares four activation implementations under an unrecorded benchmark configuration. The chart preserves the original relative illustration but does not establish portable performance ratios. The y-axis is truncated to begin at 0.40 s, so bar heights exaggerate the differences.

Modern accelerators can implement these functions through native instructions, polynomial approximations, lookup methods, or fused kernels. ReLU requires a simple threshold operation, while sigmoid, tanh, and exact Gaussian Error Linear Unit (GELU) formulations require more arithmetic or special-function support. The realized latency difference still depends on whether arithmetic, memory traffic, or launch overhead is binding. ReLU may produce zeros that a later system can exploit, but the sparsity fraction depends on the activation distribution and trained

model. Softmax¹⁰ adds a reduction and normalization across the selected axis. Those dependencies require coordination, though optimized implementations remain highly parallel.

Table 8.3 summarizes how these behaviors translate into system constraints.

Table 8.3: Activation Function Systems Comparison: Activation functions differ in mathematical behavior, kernel cost, fusion opportunities, and memory access patterns.

Function	Key Advantages	Key Disadvantages	System Implications
Sigmoid	Smooth; bounded output in (0, 1).	Gradients saturate at large magnitudes; output is not zero-centered.	Uses exponential or approximation kernels whose cost depends on the target.
Tanh	Smooth; zero-centered output in (−1, 1).	Gradients saturate at large magnitudes.	Uses special-function or approximation kernels similar in structure to sigmoid.
ReLU	Simple threshold operation; nonsaturating for positive inputs.	Can produce persistently inactive units.	Often inexpensive and fusible; exploitable sparsity depends on downstream support.
Softmax	Produces a normalized distribution over an axis.	Requires reductions and cross-element coordination.	Parallel implementations coordinate maximum, sum, and normalization stages.

ReLU remains common because it is simple and efficient, while softmax is essential when a model must normalize scores into probabilities. GPT-2 illustrates a different choice through GELU, the Gaussian Error Linear Unit.

🔍 Systems Perspective 8.3: The GELU activation choice

Beyond the foundational activation functions covered in chapter 5 (Sigmoid, Tanh, ReLU), modern architectures increasingly adopt smoother alternatives. GPT-2 uses a GELU activation (Radford et al. 2019), the Gaussian Error Linear Unit introduced by Hendrycks and Gimpel (2016) and often implemented with the common tanh-based approximation from that formulation. Its exact definition is:

$$\text{GELU}(x) = x \cdot \Phi(x) = x \cdot \frac{1}{2} \left[1 + \text{erf} \left(\frac{x}{\sqrt{2}} \right) \right]$$

where $\Phi(x)$ is the cumulative distribution function of the standard normal distribution. GELU applies a smooth, input-dependent scaling rather than ReLU’s hard threshold. It is deterministic despite its formulation in terms of a Gaussian cumulative distribution. Whether it improves model quality relative to another activation is an empirical property of the architecture and training recipe, not a universal guarantee.

Exact GELU uses `erf`, while common implementations use a tanh-based approximation such as listing 8.1. The approximation changes arithmetic cost and introduces a small numerical difference from the exact function. End-to-end impact depends on fusion, tensor size, precision, kernel implementation, and the fraction of runtime spent in activation kernels, so no fixed GELU-to-ReLU ratio applies across systems.

Listing 8.1: GELU Approximation: Fast approximation avoids expensive `erf()` computation while preserving activation properties.

```
# Fast GELU approximation used in production systems
# Avoids expensive erf() computation while
# preserving activation properties
gelu_approx = (
    0.5 * x * (1 + tanh(sqrt(2 / pi) * (x + 0.044715 * x**3)))
)
```

¹⁰ | **Softmax:** Computing a numerically stable softmax requires a maximum reduction, exponentiation, a sum reduction, and normalization across the selected axis. Parallel kernels implement these stages with coordinated reductions. FlashAttention restructures the larger attention computation so softmax statistics can be maintained while tiles remain near the compute units.

The GELU approximation highlights a broader pattern: compute cost is not always the dominant concern. For activation functions, the real bottleneck is often memory bandwidth rather than arithmetic operations. This distinction between compute-bound and memory-bound operations directly affects optimization priorities and recurs throughout the analysis of training bottlenecks.

Dense matrix multiplication can reuse operands across many multiply-accumulate operations, while an unfused elementwise activation typically performs little work for each value read and written. This arithmetic-intensity gap often makes standalone activation kernels memory-sensitive. Fusion can remove intermediate memory traffic, and special-function throughput can still matter for kernels whose arithmetic is not fully hidden. Replacing one activation with another therefore has no fixed wall-clock effect.

Systems Perspective 8.4: Memory bandwidth bottlenecks

Activation functions reveal a critical systems principle: not all operations are compute bound. While matrix multiplications saturate accelerator compute units, activation functions often become memory bandwidth bound for three reasons:

- Elementwise operations perform few calculations per memory access; ReLU performs one operation per load.
- Simple operations complete faster than memory transfer time, limiting parallelism benefits.
- Accelerator compute capability has grown faster than external-memory bandwidth, increasing pressure to reuse data near the compute units.

This is why scalar operation counts alone do not predict activation-kernel latency. The forward pass must account for fusion and activation storage as well as function-evaluation cost when determining whether memory bandwidth limits throughput.

Forward pass operations and their computational characteristics establish the workload that training systems must compute: matrix multiplications dominating FLOPs, activation functions constrained by memory bandwidth. A neural network that only computes predictions, however, learns nothing. Training requires updating model parameters so future predictions improve. The forward pass produces a loss value quantifying how wrong the current predictions are; the question now shifts from *how* much computation costs to *how* to use the result to improve.

8.3.2 Optimization algorithms

Optimization algorithms determine, given a loss value and the gradient information it produces, how each parameter should change to reduce future errors. These algorithms govern the learning trajectory, translating gradients into parameter updates that steer the model toward better performance, and their selection has direct system-level implications for computation efficiency, memory requirements, and scalability. The focus here is on the algorithms themselves during training: how they use gradients, how much state they retain, and how their update rules interact with the hardware budget.

8.3.2.1 Gradient-based optimization methods

Section 5.4.4.5 introduces gradient descent as the fundamental optimization algorithm: iteratively adjusting parameters in the direction of steepest descent. That conceptual foundation assumed modest networks on single devices. At hardware scale, gradient descent and its variants interact with real physical constraints. The same mathematical operation that elegantly adjusts weights becomes a significant systems challenge when models contain billions of parameters and training data spans terabytes.

Gradient descent Gradient descent is the mathematical foundation of neural network training, iteratively adjusting parameters to minimize a loss function. In training systems, this mathematical operation translates into specific computational patterns. For each iteration, the system must execute four dependent operations:

1. Compute forward pass activations
2. Calculate loss value
3. Compute gradients through backpropagation
4. Update parameters using the gradient values

The computational demands of gradient descent scale with both model size and dataset size. Computing gradients requires storing intermediate activations during the forward pass for backpropagation. These activations consume memory proportional to the depth of the network and the number of examples being processed.

Traditional gradient descent processes the entire dataset before each parameter update. For a training set with one million examples, the system must compute an aggregate gradient over all examples before taking one step. Let D denote the number of examples in the training dataset, let B_{micro} denote the examples resident for one forward/backward pass, and let $T_{\text{forward+backward per example}}$ denote the amortized forward/backward time per example under the chosen micro-batch execution regime. Gradient accumulation composes multiple micro-batches into one larger effective batch. The peak-memory relation in equation 8.2 and the step-time relation in equation 8.3 capture the implementation distinction:

$$\text{Peak Activation Memory} \propto B_{\text{micro}} \times \text{Activation Memory per Example} \quad (8.2)$$

$$T_{\text{step}} \propto D \times T_{\text{forward+backward per example}} \quad (8.3)$$

Section C.3.3 formalizes this breakdown by deriving the full training memory equation, including optimizer-state overhead. Full-batch training does not require storing every example's activations simultaneously if the dataset is streamed or microbatched; peak memory is governed by the examples held in memory at once. The systems problem is still severe: processing $D = 1,000,000$ examples before each update creates million-example iteration times, reducing the rate at which the model can learn from the data.

Hardware-friendly variants relax the need for an exact full-dataset gradient on every update. Stochastic gradient descent (SGD)¹¹ estimates gradients from sampled examples or mini-batches rather than the entire dataset. This changes update frequency and gradient variance while allowing each update to operate on a bounded working set.

However, processing single examples creates new system challenges. Modern accelerators achieve peak performance through parallel computation, processing multiple data elements simultaneously. Single-example updates leave most computing resources idle, resulting in poor hardware utilization. The frequent parameter updates also increase memory bandwidth requirements, as weights must be read and written for each example rather than amortizing these operations across multiple examples.

Mini-batch processing Mini-batch gradient descent emerges as a practical compromise between full-batch and stochastic methods, an algorithm-machine co-design that computes gradients over small batches of examples aligned with modern accelerator architectures (Dean et al. 2012). The batch size B becomes a key system parameter, influencing both computational efficiency and memory requirements.



Definition 8.3: Batch processing

Batch processing aggregates multiple training examples into tensor operations that amortize fixed per-step overhead across B examples and can expose more parallel work. Increasing B may improve occupancy and arithmetic intensity, but it does not guarantee that every operation becomes compute bound.

1. Significance: Throughput often increases with batch size until the target hardware is sufficiently occupied, while the number and quality of updates needed for convergence follow a separate workload-dependent curve. Goyal et al. demonstrated a ResNet-50 ImageNet recipe at $B = 8,192$ using linear learning-rate scaling and warmup (Goyal et al. 2017). That result is a configuration-specific demonstration rather than a universal critical batch size.

¹¹ | **Stochastic gradient descent:** “Stochastic” from Greek *stochastikos* (“able to guess”); rather than waiting for an exact full-dataset gradient, SGD updates from sampled examples or mini-batches. Full-batch and stochastic methods can both stream data through bounded micro-batches, so their peak activation memory is determined by the resident micro-batch rather than by dataset size alone. Their central difference is how many sampled gradients contribute to each parameter update.

2. Distinction: A mini-batch update averages gradients over B examples. Under independent, similarly distributed samples, the standard deviation of this mean decreases approximately as $1/\sqrt{B}$. The update processes more data than a single-example update, so its net data movement and time depend on reuse, implementation, and how batch size changes the number of updates required.
3. Common pitfall: A frequent misconception is that linear learning-rate scaling (multiplying η by B/B_0) works at any batch size. It is a training recipe that commonly requires warmup and empirical validation, and its useful range depends on the model, data, optimizer, and target accuracy.

The relationship between batch size and system performance reveals hardware-software trade-offs. The decomposition in equation 8.4 separates fixed parameter and gradient storage from a batch-dependent activation term:

$$\text{Memory Required} = \text{Parameter Memory} + \text{Gradient Memory} + B \times \text{Activation Memory} \quad (8.4)$$

Because the activation term scales with B while parameter and gradient memory stay fixed, doubling the batch doubles the activation working set, and a model that fits comfortably at a small batch can exhaust the 40–80 GB of HBM on a high-end training accelerator once the batch grows. Section 8.4.3.2 works this budget through layer by layer for ResNet-50.

Larger batches can enable more efficient computation through improved parallelism and better memory access patterns. Accelerator utilization efficiency demonstrates this trade-off: larger batches generally expose more parallel work to the accelerator, while very small batches can leave compute units underfilled. Linear scaling rules for large-batch training (scale learning rate proportionally to batch size increase) can help maintain convergence speed within an empirically validated range (Goyal et al. 2017).

This establishes a central theme in training systems: the hardware-software trade-off between memory constraints and computational efficiency. Training systems must select batch sizes that maximize hardware utilization while fitting within available memory. The optimal choice often requires gradient accumulation when memory constraints prevent using efficiently large batches, trading micro-batch serialization and some overhead for the same effective batch size.

8.3.2.2 Adaptive and momentum-based optimizers

Basic SGD applies one learning rate to every parameter, which can make optimization difficult when gradient scales differ substantially across dimensions. Momentum smooths updates using gradient history, RMSprop adapts step sizes per parameter, and Adam combines first-moment tracking with adaptive scaling (Kingma and Ba 2015). These methods consume gradients computed by backpropagation; they do not determine whether those gradients are correct. Their convergence behavior depends on the model and training recipe, while their additional state creates concrete memory and computation costs.

Momentum-based methods Momentum methods¹² address SGD’s oscillation problem by accumulating a velocity vector across iterations, smoothing out noisy gradient directions. From a systems perspective, this smoothing comes at a cost: the training system must maintain a velocity vector with the same dimensionality as the parameter vector. Parameters plus this auxiliary state therefore occupy twice the parameter storage before gradients and activations are counted.

Adaptive learning rate methods While momentum smooths gradient direction, it does not address the different scales of gradients across parameters. RMSprop¹³ solves this by maintaining a moving average of squared gradients for each parameter, automatically reducing step sizes for parameters with historically large gradients. This per-parameter adaptation requires storing the moving average s_t , creating memory overhead similar to momentum methods. The elementwise operations in RMSprop also introduce additional computational steps compared to basic gradient descent.

Adam optimization Adam¹⁴ combines the benefits of both momentum and RMSprop: momentum’s gradient smoothing addresses noisy updates, while RMSprop’s adaptive scaling handles parameter-specific step sizes. This combination maintains two moving averages for each parameter. Here

¹² | **Momentum:** Borrowed from physics, where momentum (mass $\times$ velocity) describes an object’s tendency to continue moving. The metaphor explains the design because accumulated velocity smooths noisy gradients and can reduce the iterations needed for convergence. The systems cost is one additional velocity value per parameter, while Adam maintains two moment values per parameter.

¹³ | **RMSprop:** Proposed by Geoffrey Hinton in Lecture 6e of his 2012 Coursera course—never published in a peer-reviewed paper, making it perhaps the most influential optimizer disseminated via a slide deck. RMSprop divides the learning rate by a running average of recent gradient magnitudes, adapting step sizes per parameter. This per-parameter adaptation is what Adam inherits as its second moment v_t , directly contributing to Adam’s $3\times$ memory overhead described in section 8.3.2.3.

¹⁴ | **Adam (adaptive moment estimation):** The two moving averages are the first moment (momentum) and second moment (uncentered variance) of the gradients, stored for every model parameter. For a 7B model, Adam’s FP32 moment tensors alone consume 56 GB; the full standard mixed-precision training state before activations is 112 GB once FP16 weights, FP16 gradients, FP32 master weights, and Adam moments are counted together.

m_t and v_t are the first- and second-moment buffers, β_1 and β_2 are their decay factors, and ϵ is the numerical-stability constant in the denominator:

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) \nabla \mathcal{L}(\theta_t) \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) (\nabla \mathcal{L}(\theta_t))^2 \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t} \\ \theta_{t+1} &= \theta_t - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon} \end{aligned}$$

The system implications of Adam are more substantial than previous methods. The optimizer must store two additional vectors (m_t and v_t) for each parameter, tripling the parameter-plus-optimizer-state footprint; for a 100M-parameter model, those auxiliary vectors alone add 800 MB beyond weight storage.

8.3.2.3 Optimization algorithm system implications

The choice of optimization algorithm creates specific patterns of computation and memory access that influence training efficiency. Optimizer auxiliary memory increases progressively from SGD (no auxiliary state) through Momentum (one velocity vector) to Adam (two moment vectors), as quantified in table 8.4. These memory costs must be balanced against convergence¹⁵ benefits. Adam may require fewer iterations for some workloads, but its per-iteration memory and computation overhead can reduce training speed on memory-constrained systems. At GPT-2 scale, this overhead becomes a first-order memory constraint.

Table 8.4: Optimizer Memory Footprint: Different optimization algorithms impose varying auxiliary-state costs due to the storage of intermediate values like velocities and squared gradients. The multiplier row counts parameters plus optimizer auxiliary state and excludes gradients and activations; full training memory must add those terms explicitly. These trade-offs govern resource-constrained deployments and large-scale model training.

Property	SGD	Momentum	RMSprop	Adam
Memory Overhead	None	Velocity terms	Squared gradients	Both velocity and squared gradients
Parameter + optimizer state	1×	2×	2×	3×
Dense State Access	Parameters and gradients	Adds velocity	Adds squared-gradient state	Adds first- and second-moment state
Update Work	Lowest auxiliary work	Adds momentum update	Adds adaptive scaling	Adds moment updates and bias correction
Hardware Efficiency	Varies	Varies	Varies	Varies
Convergence Behavior	Varies	Varies	Varies	Varies

The costs quantified in table 8.4 create a design tension. Adam’s two moment buffers may improve optimization for a given workload, but their memory can reduce the model size or batch size that fits on an accelerator. AdamW¹⁶ (Loshchilov and Hutter 2019) decouples weight decay from Adam’s adaptive gradient update without adding another per-parameter state tensor.

 Napkin Math 8.2: GPT-2 optimizer memory requirements

A representative GPT-2 XL training configuration uses the Adam optimizer with five hyperparameters:

- $\beta_1 = 0.9$ (momentum decay)
- $\beta_2 = 0.999$ (second moment decay)
- Learning rate: Warmed up from 0 to 2.5e-4 over first 500 steps, then cosine decay
- Weight decay: 0.01
- Gradient clipping: Global norm clipping at 1.0

¹⁵ | **Convergence:** A run reaches its stopping criterion when the selected training or validation metric no longer improves enough to justify additional work. The required steps depend on the model, data, optimizer, schedule, batch size, and target quality. Fewer steps can reduce wall-clock time and cost, but adaptive optimizers add state whose memory and traffic must be included in the system budget.

¹⁶ | **AdamW (Adam with decoupled weight decay):** Standard L_2 regularization enters the gradient and is therefore rescaled by Adam’s adaptive update. AdamW applies weight decay as a separate parameter update, restoring the intended distinction between adaptive gradient scaling and decay (Loshchilov and Hutter 2019). It uses the same moment-buffer structure as Adam, though its effect on model quality remains workload dependent.

Math:

For GPT-2's 1.5B parameters in FP32 (4 bytes each), the memory breaks down across four components:

- Parameters: $1.5\text{B} \times 4 \text{ bytes} = 6 \text{ GB}$
- Gradients: $1.5\text{B} \times 4 \text{ bytes} = 6 \text{ GB}$
- Adam State (m, v): $1.5\text{B} \times 8 \text{ bytes} = 12 \text{ GB}$
- Total static memory: 24 GB

This explains why GPT-2's 24 GB static training state alone approaches the 32 GB capacity of a V100 before activation storage is counted.

System Decisions Driven by Optimizer

1. Mixed precision training (FP16) reduces operation precision but requires keeping FP32 master weights, maintaining the static memory footprint at ~24 GB.
2. Gradient accumulation (splitting effective batches into smaller micro-batches) allows effective batch size $B = 512$ despite memory limits.

Systems insight: This illustrative configuration shows how Adam's moment buffers constrain capacity. Any convergence benefit relative to SGD or momentum must be measured under a matched model, data, schedule, batch size, and stopping criterion.

8.3.2.4 Framework optimizer interface and scheduling

After optimizer memory is quantified, the framework interface matters because it fixes when that state is read, written, cleared, and preserved across steps. A training loop separates gradient computation from parameter updates so that the system can accumulate gradients, synchronize them, or defer updates without changing the optimizer equations. Listing 8.2 demonstrates where Adam optimization enters that cycle.

Listing 8.2: Adam Training Loop: Standard four-step optimization cycle with gradient clearing, forward pass, backward pass, and parameter update.

```
import torch
import torch.nn as nn
import torch.optim as optim

# Initialize Adam optimizer with model parameters and learning rate
optimizer = optim.Adam(
    model.parameters(), lr=0.001, betas=(0.9, 0.999)
)
loss_function = nn.CrossEntropyLoss()

# Standard training loop implementing the four-step optimization cycle
for epoch in range(num_epochs):
    for batch_idx, (data, targets) in enumerate(dataloader):
        # Step 1: Clear accumulated gradients from previous iteration
        optimizer.zero_grad()

        # Step 2: Forward pass - compute model predictions
        predictions = model(data)
        loss = loss_function(predictions, targets)

        # Step 3: Backward pass - compute gradients via autodiff
        loss.backward()

        # Step 4: Parameter update - apply Adam optimization equations
        optimizer.step()
```

The `optimizer.zero_grad()` call marks the boundary between one update and the next. Gradients accumulate across calls to `backward()`, so clearing them explicitly prevents stale gradients from contaminating the next batch. The same accumulation behavior later becomes useful for large effective batch sizes, but only when the training loop manages the boundary deliberately.

The `optimizer.step()` method is the other boundary: it consumes the current gradients and mutates persistent optimizer state. For Adam optimization, this call implements momentum estimation, squared gradient tracking, bias correction, and the parameter update. Algorithm 8.1 makes the hidden state explicit so the memory cost remains visible rather than disappearing behind an API call.

Algorithm 8.1 Adam parameter update (one optimizer step)

Require: gradient $g_t = \nabla \mathcal{L}(\theta_t)$; step t ; rate η ; decays β_1, β_2 ; constant ϵ

Ensure: updated parameter θ_{t+1} ; moment buffers m_t, v_t carried to the next step

- 1: $m_t \leftarrow \beta_1 m_{t-1} + (1 - \beta_1) g_t \triangleright$ first moment (momentum)
 - 2: $v_t \leftarrow \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \triangleright$ uncentered second moment
 - 3: $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$; $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t) \triangleright$ bias correction
 - 4: $\theta_{t+1} \leftarrow \theta_t - \eta \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon) \triangleright$ parameter update
-

P	G	Adam
---	---	------

In this simplified static-state view, Adam's two moment buffers occupy twice the parameter storage.

Steps 1 and 2 keep two persistent moment buffers per parameter, so parameters plus optimizer state reach roughly $3 \times$ the parameter memory before gradients, activations, and FP32 master weights are counted. Framework implementations manage the allocator and access patterns for these optimizer states, but they do not remove the cost. Each Adam step reads the gradient, parameter, first moment, and second moment, then writes back the updated moments and parameter. The abstraction reduces implementation burden; the systems budget still pays for two extra tensors that must occupy memory and move through the hierarchy every step.

8.3.2.5 Learning rate scheduling integration

The framework's learning rate scheduling hook changes the optimizer's trajectory without adding per-parameter state. It adjusts the learning rate η during training, letting the system shape convergence behavior while leaving the underlying optimizer equations intact.

Schedules such as cosine annealing, exponential decay, or step-wise reductions implement that trajectory by changing the step size over time. A scheduler may track the current step, epoch, warmup phase, or other state that must be restored with a checkpoint. It changes η without changing the optimizer's base update equations; chapter 7 covers the scheduler interface that wires this in. This separation lets the systems engineer combine base optimization algorithms (SGD, Adam) with scheduling strategies (cosine annealing, linear warmup) without reimplementing the update rule.

The optimization algorithms in the preceding section specify how to update parameters given gradients, but they take those gradients as given. SGD, momentum, and Adam all assume gradient vectors arrive ready-made. In practice, computing gradients for a network with billions of parameters is itself a major computational and memory challenge. The cost of gradient computation, not the cost of the optimizer step, is what makes training so much more expensive than inference.

8.3.3 Backpropagation mechanics

Backpropagation solves the gradient computation problem by tracing error signals backward through the network, systematically attributing responsibility to each parameter for the final prediction error. Its memory and computational requirements reveal why training systems face such substantial resource constraints.

The backpropagation algorithm computes gradients by systematically moving backward through a neural network's computational graph. Section 5.4.4 establishes the mathematical foundation: the chain rule breaks gradient computation into layer-by-layer operations, with each layer receiving adjustment signals proportional to its contribution to the final error. If terms like "computational graph" or "gradient flow" feel unfamiliar, the factory assembly line analogy in that section is worth revisiting.

At system scale, the question shifts from *what* backpropagation computes to *what* it costs. The layer computations from section 8.3.1.1 produce activations that must be retained for the backward pass. Computing $\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(\ell)}}$ requires access to these stored activations, creating the training memory equation that gradient checkpointing later exploits.

A simple three-layer network processing MNIST requires kilobytes of activation storage. GPT-2 processing a single batch requires over 35.9 GB, more than most accelerators can hold. That gap defines the engineering challenge this chapter addresses. Section C.3.3 derives how backpropagation drives these memory costs, including the full training memory equation ($M_{\text{total}} = M_{\text{weights}} + M_{\text{gradients}} + M_{\text{optimizer}} + M_{\text{activations}}$). Modern training systems use autodifferentiation (see chapter 7) to handle gradient computations automatically, but the underlying memory and computation patterns remain the systems engineer's responsibility to manage.

8.3.3.1 Activation memory requirements

Training systems retain intermediate values requested by backward rules. The exact saved set depends on the operation, fusion, and recomputation strategy. Training state includes several distinct categories:

- saved forward values needed by backward rules,
- model parameters,
- parameter gradients, and
- optimizer state and any higher-precision parameter copies.

Consider a batch of training examples passing through a network. The forward pass computes and stores:

$$\mathbf{Z}^{(\ell)} = \mathbf{A}^{(\ell-1)} \mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)}$$

$$\mathbf{A}^{(\ell)} = f(\mathbf{Z}^{(\ell)})$$

A conventional implementation may retain $\mathbf{Z}^{(\ell)}$, $\mathbf{A}^{(\ell)}$, or other inputs and outputs according to the derivative rule. Fused kernels and recomputation can retain a smaller set. Batch-dependent activations still create multiplicative pressure, while optimizer overhead scales with parameter count. The GPT-2 calculation decomposes that pressure into per-layer attention, feed-forward, and training-state costs.

Napkin Math 8.3: GPT-2 activation memory breakdown

For GPT-2 with batch size $B = 4$, sequence length $S = 1024$, hidden dimension $d = 1600$, and 48 layers:

Math:

- Attention activations: $B \times S \times d \times 4$ (Q, K, V, output) = $4 \times 1024 \times 1600 \times 4 \times 2$ bytes (FP16) = 52.4 MB
- FFN activations: $B \times S \times (4d)$ (intermediate expansion) = $4 \times 1024 \times 6400 \times 2$ bytes = 52.4 MB
- Attention scores: $5 \times N_{\text{heads}} \times S^2 \times B$ bytes for the $S \times S$ score, softmax, and dropout buffers. With 25 heads, this term reaches 524.3 MB per layer—the dominant term, quadratic in sequence length, and exactly what selective recomputation discards
- Layer norm states: Minimal (~10 MB per layer)
- Total per layer: ~639.1 MB (attention + FFN + attention scores + layer norm states)

Result:

- Total activation memory (library estimate, includes residual-stream and framework buffers beyond the line items above): 35.9 GB
- Parameters (FP16): 3 GB
- Gradients: 3 GB

GPT-2 acts 71.7 GB

V100 32 GiB

MNIST 438 KB

log scale

Activation memory spans
MNIST toys to GPT-scale
training.

- Master weights (FP32): 6 GB
- Adam moment state (FP32): 12 GB
- Peak memory during training: ~59.9 GB

This exceeds a single V100's 32 GB capacity.

Solutions:

1. Selective activation recomputation: Discard attention intermediates and recompute them during backward, reducing this modeled activation term by 70.2 percent to ~10.7 GB; the additional work depends on the implementation
2. Activation CPU offloading: Store some activations in CPU RAM, transfer during backward pass
3. Mixed precision: FP16 activations (already applied) vs. FP32 (would be 71.7 GB)
4. Reduced micro-batch size and gradient accumulation: Build a larger effective batch from multiple memory-feasible forward/backward passes

Systems insight: Even after selective recomputation, the modeled activations plus static state total about 34.7 GB, slightly above the raw 32 GB V100 capacity before workspace and allocator overhead. Full recomputation, sharding, offload, or a smaller micro-batch is therefore still required for this representation.

This breakdown illustrates the practical engineering decisions required when accelerator memory falls short. The same trade-off between stored activations, recomputation, and batch size drives the memory-computation analysis that follows.



Checkpoint 8.2: The memory-compute trade-off

Training large models requires managing the memory wall (the bandwidth bottleneck introduced in chapter 5 and revisited in section 7.3.1).

Bottleneck

- ☐ Activation memory: Do you understand why activations stored for backpropagation can dominate memory usage and, for some model shapes and training recipes, exceed parameter size by more than 10×?
- ☐ Optimization strategy: Can you explain how gradient checkpointing trades compute (re-calculating activations) for memory capacity?
- ☐ **Memory estimate:** for GPT-2 (1.5B parameters) in mixed precision, put a one-significant-figure number on each term in equation 8.6: weights, gradients, optimizer state, and activations at batch size 32. Which term is largest, and would the total fit a 32 GB accelerator?

Scaling limits

- ☐ **Batch size constraints:** why does memory capacity limit the maximum batch size, and how does gradient accumulation solve this without increasing memory?

8.3.3.2 Memory-computation trade-offs

Training systems must balance memory usage against computational efficiency. Each forward pass generates values that backward rules may request. For a simplified network model with N_L layers, let s_ℓ denote the bytes of saved intermediate state per example at layer ℓ and a_ℓ the bytes of saved activation outputs per example. Under this per-example convention, their storage scales linearly with batch size and is approximated by equation 8.5:

$$\text{Memory per batch} = B \times \sum_{\ell=1}^{N_L} (s_\ell + a_\ell) \quad (8.5)$$

This memory requirement compounds with the weights, gradients, and optimizer memory discussed in section 8.3.2.3. Equation 8.6 gives a conceptual training-state decomposition:

$$\text{Total Memory} = \text{Memory}_{\text{weights}} + \text{Memory}_{\text{gradients}} + \text{Memory}_{\text{optimizer}} + \text{Memory per batch} \quad (8.6)$$

The runtime peak can exceed this sum because temporary workspaces, communication buffers, allocator behavior, and framework bookkeeping also occupy memory. To reduce the activation term, training systems can strategically recompute intermediate values during backward rather than storing them. This increases computational work but can enable deeper networks or larger batches on memory-constrained hardware (T. Chen et al. 2016). Algorithm 8.2 makes the trade explicit: the forward pass keeps activations at only a sparse set of checkpoint layers, and the backward pass recomputes the rest on demand.

Algorithm 8.2 Gradient checkpointing (activation recomputation)

Require: N_L -layer network; checkpoint set $\mathcal{C} \subseteq \{1, \dots, N_L\}$ (e.g., every $\sqrt{N_L}$ layers)

Ensure: parameter gradients, at reduced peak activation memory

```

1: for  $\ell = 1$  to  $N_L$  do
2:   forward layer  $\ell$ ; store its activation only if  $\ell \in \mathcal{C} \triangleright$  drop the rest
3: end for
4: for  $\ell = N_L$  down to 1 do
5:   if the activation of layer  $\ell$  was dropped then
6:     recompute the forward segment from the nearest stored checkpoint through layer  $\ell$ 
7:   end if
8:   compute the layer's gradient; free the activation
9: end for
10: return the accumulated gradients

```

If we space checkpoints every $\sqrt{N_L}$ layers, we can reduce peak activation memory from $\mathcal{O}(N_L)$ to $\mathcal{O}(\sqrt{N_L})$ by doing extra forward work over checkpoint segments during the backward pass. That is the lever that fits a network too large to train within accelerator memory, trading the iron law's operation term for activation capacity.

The efficiency of these memory management strategies depends heavily on the underlying hardware architecture. Accelerator systems, with their high computational throughput but limited memory bandwidth, often encounter different bottlenecks than CPU systems. Memory bandwidth limitations on accelerators mean that even when sufficient storage exists, moving data between memory and compute units can become the primary performance constraint (Jouppi et al. 2017).

These hardware considerations guide the implementation of backpropagation in modern training systems. Specialized memory-efficient algorithms for operations like convolutions compute gradients in tiles or chunks, adapting to available memory bandwidth. Dynamic memory management tracks the lifetime of intermediate values throughout the computation graph, deallocating memory as soon as tensors become unnecessary for subsequent computations (Paszke et al. 2019).

Forward propagation, gradient computation, and parameter updates define *what* training systems must compute. An operation's name alone, however, does not reveal *where* the system stalls. Matrix shape, fusion, precision, implementation, and hardware determine whether computation or data movement is binding. Arithmetic intensity provides the next analytical tool.

8.3.4 Arithmetic intensity

Arithmetic intensity captures this distinction—the ratio of computation to data movement that reveals whether an operation is limited by compute throughput or memory bandwidth:

$$\text{Arithmetic Intensity} = \frac{\text{FLOPs}}{\text{bytes moved}}$$

Operations with high arithmetic intensity are compute bound: their performance is limited by the processor's computational throughput. Operations with low arithmetic intensity are memory

bound: they spend more time moving data than computing. Section D.2.1 gives the formal definition of the Roofline Model and shows how to compute a hardware’s ridge point.¹⁷

Table 8.5 summarizes common tendencies rather than fixed measurements. Dense matrix multiplication can reuse operands enough to reach high arithmetic intensity, while standalone elementwise and reduction kernels often perform less work per byte moved. The target hardware’s ridge point determines the resulting regime.

Table 8.5: Training Operation Classification Conditions: Operation families exhibit common arithmetic-intensity tendencies, but classification requires comparing the measured implementation with the target hardware’s ridge point. Memory-sensitive kernels may benefit from reducing traffic or increasing fusion, while compute-bound kernels benefit from additional arithmetic throughput.

Operation	Arithmetic-intensity tendency	Classification condition
Dense MatMul (large)	High when shapes and tiling provide substantial operand reuse	Compute-bound only when intensity exceeds the ridge
Activation functions	Usually low as standalone elementwise kernels	Often memory-sensitive unless fusion removes traffic
LayerNorm/BatchNorm	Reduction and elementwise work with multiple tensor passes	Depends on fusion, shape, and target ridge point
Attention softmax	Reduction and normalization over the attention axis	Depends on materialization, tiling, and fusion

Hardware efficiency during model training depends directly on kernel arithmetic intensity. Locate the operation points in figure 8.2 along the logarithmic arithmetic intensity axis, comparing memory-bandwidth bound operators on the sloped ceiling against compute-bound operators on the flat peak.




Figure 8.2: Illustrative Training Roofline: An A100 roofline with assumed operation coordinates illustrates how arithmetic intensity changes the attainable-performance ceiling. The points are not a profile of one GPT-2 execution. FlashAttention is shown moving right because it reduces HBM traffic; whether a real attention kernel crosses the ridge depends on shape, precision, hardware, and implementation.

To build intuition for these relationships, study the roofline diagram in figure 8.2. The ridge point marks the “knee” where the sloped memory-bound region meets the flat compute-bound ceiling. An implementation plotted left of this point is bandwidth bound under the roofline assumptions; one plotted right is compute bound. The operation points in this diagram are illustrative placements that show how to interpret the model, not measured GPT-2 coordinates.

Consider a GPT-2 attention layer that materializes the $S \times S$ attention-score matrix. For each head, the $\mathbf{QK}^\top$ product costs approximately $2S^2d_{\text{head}}$ FLOPs, while writing the FP16 score matrix and reading it back moves about $4S^2$ bytes. Counting only that score-matrix traffic gives the

¹⁷ | **Ridge point and precision:** The roofline ridge point—the arithmetic intensity threshold separating memory-bound from compute-bound operations—shifts with numerical precision. On the same accelerator, a lower-precision Tensor Core path can expose much more arithmetic throughput at roughly the same memory bandwidth, raising the ridge point substantially. Switching from TF32-style execution to BF16 mixed precision can therefore change which optimization technique yields returns, making precision selection inseparable from roofline analysis.

approximation $d_{\text{head}}/2$. For GPT-2 Small ($d_{\text{model}} = 768$ across 12 heads, so $d_{\text{head}} = 64$), this restricted accounting yields 32 FLOP/byte. A complete kernel analysis must also count Q and K traffic, softmax work, additional reads and writes, cache reuse, and fusion. The $\mathcal{O}(S^2)$ materialized score traffic remains the term that FlashAttention targets.

Accelerators have characteristic hardware ridge points where operations transition from memory-bound to compute bound. A representative data-center accelerator has a ridge point high enough that low-intensity operations such as materialized attention softmax remain memory bound even when large matrix multiplications are compute bound. Operations below the ridge point are memory bound; above it, they are compute bound.



Systems Perspective 8.5: Peak FLOP/s vs. sustained performance

Peak TFLOP/s is an arithmetic ceiling rather than a prediction of sustained application performance. For an implementation that is strictly bandwidth bound, raising only the arithmetic ceiling does not improve the roofline bound. Mixed-precision training can both enable faster arithmetic paths and reduce traffic for tensors actually stored and transferred at lower precision. The realized gain depends on kernel support, conversion overhead, numerical requirements, and which resource is binding. Roofline-guided optimization therefore measures the implementation before choosing fusion, precision changes, or additional compute.

Batch size can influence arithmetic intensity and occupancy by changing matrix shapes and the amount of reuse available to a kernel. No universal boundary at batch 32 separates memory-bound from compute-bound execution; the transition depends on the operation, dimensions, implementation, precision, and hardware.

This analysis guides optimization strategy selection. For memory-bound operations, reducing data movement through operator fusion, reduced precision, or algorithmic improvements like FlashAttention provides the largest gains. For compute-bound operations, increasing throughput through Tensor Cores and parallel execution matters more. The distinction is practical: the first case asks how to move fewer bytes, while the second asks how to keep more arithmetic units busy.

In figure 8.2, FlashAttention¹⁸ captures the core insight of IO-aware algorithm design. By never materializing the full $S \times S$ attention matrix in HBM and instead processing tiles that fit in fast SRAM, FlashAttention reduces auxiliary attention memory from $\mathcal{O}(S^2)$ to $\mathcal{O}(S)$ and sharply reduces HBM traffic. The original work reports up to $3\times$ speedups on evaluated workloads (T. Dao et al. 2022), with gains depending on tensor shape, hardware, and baseline. The algorithm and the conditions under which it applies are examined in detail in section 8.6.4.

The arithmetic-intensity analysis in this section shows how to determine which resource constrains a particular implementation. Dense matrix multiplication often reaches the compute-bound regime when shapes provide enough reuse, while standalone normalization and activation kernels are often memory-sensitive. FlashAttention exemplifies how an algorithm can reduce data movement and raise arithmetic intensity, though the final regime remains workload dependent.

Optimizing individual operations is necessary but insufficient. A perfectly tuned matrix multiplication achieves nothing if the accelerator sits idle waiting for the next batch of data. The mathematical foundations established earlier in this chapter quantified the cost of each piece—matrix multiplications consuming trillions of FLOPs, activation functions bottlenecked by memory bandwidth, optimizer states tripling memory requirements. The next question is *how* to orchestrate these pieces into a pipeline where no stage starves the others.

8.4 Pipeline Architecture

A training step is not a single operation but a sequence of dependent stages—data must be loaded before computation can begin, forward passes must complete before backward passes start, and gradients must be computed before parameters can update. The speed of the slowest stage determines the speed of the entire system.

The system-level pipeline coordinates these stages across real hardware with finite memory and bandwidth constraints. Chapter 7 introduced how frameworks like PyTorch and TensorFlow provide APIs for defining models and executing forward passes; here those API calls become part of a larger

¹⁸ | **FlashAttention:** The core mechanism processes attention in small tiles that fit within the accelerator's fast on-chip SRAM, avoiding writes of the full intermediate $S \times S$ matrix to slower HBM. The exact HBM IO bound depends on tile size and SRAM capacity, but the practical effect is a large reduction in memory traffic and auxiliary activation storage. The original work reports workload-dependent speedups of up to $3\times$.

architecture of data loading, preprocessing, accelerator transfers, and parameter updates—a unified pipeline rather than isolated operations.

This orchestration is not a single monolithic process but rather three interconnected subsystems, each with distinct responsibilities and resource demands. Figure 8.3 traces how these subsystems connect: the data pipeline handles ingestion and preprocessing, the training loop executes forward passes, backward passes, and parameter updates, and the evaluation pipeline periodically assesses model quality. The flow between these components is where bottlenecks emerge—the interconnection points expose the binding constraints.




Figure 8.3: Training System Overview: The training lifecycle coordinates data preparation, core forward/backward computation, and validation evaluation. The interfaces between these three pipelines represent critical synchronization boundaries where I/O delays or evaluation pauses can throttle accelerator utilization.

8.4.1 Architectural overview

A single training iteration involves three subsystems executing in sequence: a data pipeline that ingests, transforms, and batches raw data; a training loop that performs the forward pass, gradient computation, and parameter update; and an evaluation pipeline that measures model quality against held-out data. This subsystem sequence frames the chapter’s training pipeline. Understanding each subsystem’s role clarifies where performance bottlenecks arise and where system-level optimizations have their greatest impact.

A training system is easiest to reason about from its boundaries inward. The data pipeline loads raw records from storage, applies transformations such as resizing, augmentation, and normalization, and assembles the resulting examples into batches; input normalization is a long-standing training aid (LeCun et al. 2012). The evaluation pipeline sits at the other boundary. At configurable intervals, it runs held-out validation data through the current model, computes metrics such as accuracy or loss, and exposes convergence problems such as overfitting, where training loss improves while validation loss degrades. Because evaluation consumes accelerator time, its cadence trades finer-grained feedback against training throughput.

The core engine of neural network learning sequences forward predictions, backward error propagation, and parameter optimization into a closed loop. Follow the cyclic data path in figure 8.4 across the three sequential stages: forward pass, gradient calculation, and weight updates.

Each iteration executes the forward pass, loss computation, backward pass, and parameter update cycle established in section 8.3. The systems question is not *what* these operations compute (covered earlier) but *how* they interact as a pipeline, where the bottleneck in any one stage limits overall throughput.

This process repeats across batches and epochs, gradually refining the model to improve its predictive accuracy. The loop is tightly coupled to the stages around it: data preparation can overlap with computation by preprocessing the next batch while the current batch trains, while evaluation temporarily pauses gradient updates to measure validation quality. The integration minimizes idle time for system resources, but any imbalance, such as a slow data pipeline or an overly frequent evaluation schedule, propagates as reduced overall throughput.

8.4.2 Data pipeline

The architectural overview identified the data pipeline as the first component in the training system. Its efficiency directly determines whether expensive accelerator resources remain fully engaged or




Figure 8.4: The Iterative Training Loop: Each iteration sequences forward prediction, gradient computation by backpropagation, and an optimizer update. Updated parameters return to the next forward pass while the next batch remains a separate input. Data dependencies and memory residency set the step’s minimum execution time and memory floor.

sit idle waiting for data. The systems aspects of data movement and preprocessing are the focus here; the upstream data engineering practices are covered in chapter 4.

The data pipeline running on the CPU bridges raw data storage and accelerator computation. Figure 8.5 breaks down this architecture into three distinct zones.




Figure 8.5: CPU-to-GPU Data Flow: Training data transitions through storage, CPU preprocessing, and GPU execution zones. Each stage represents a potential throughput gate; disk read bandwidth, transformation latency, and PCIe host-to-device transfers must collectively sustain the GPU consumption rate to keep accelerators at peak utilization.

These zones matter because each can become the slowest stage. Storage supplies raw examples from disk, typically image files for computer vision or text files for natural language processing. CPU preprocessing then converts formats, applies resizing, normalization, or data augmentation, and batches examples into tensors the accelerator can consume.

The GPU training zone consumes those preprocessed batches across multiple accelerators for parallel computation. Format conversion, processing, and batching are therefore not housekeeping steps; they are throughput gates. If any one runs slower than the training loop, expensive accelerator resources idle while the data pipeline catches up.

8.4.2.1 Core components

The data pipeline’s throughput is ultimately limited by how fast training data can be retrieved from storage. The data engineering practices from chapter 4, including data format selection (Parquet, TFRecord, Arrow), partitioning strategies, and data locality optimization, directly impact these storage characteristics. These storage constraints propagate through the training system.

Storage throughput is bounded by the slower of two hardware constraints, expressed in equation 8.7:

$$R_{\text{storage}} = \min(BW_{\text{disk}}, BW_{\text{network}}) \quad (8.7)$$

where BW_{disk} is the physical disk bandwidth and BW_{network} represents the network bandwidth for distributed storage systems. In practice, training workloads rarely achieve this theoretical maximum because data shuffling can turn large sequential reads into smaller, less efficient accesses. Effective

storage throughput can be modeled as $R_{\text{storage,eff}} = R_{\text{storage}} \times F_{\text{access}}$, where F_{access} depends on record size, file layout, caching, and storage hardware. This penalty explains why data pipeline engineering matters: without careful layout, prefetching, and buffering, accelerators can sit idle waiting for storage.

8.4.2.2 Preprocessing

After storage, preprocessing turns raw inputs into model-ready tensors. Data pipelines from chapter 4 commonly extract and load raw data before transforming it on demand during training. Under ideal independent-worker scaling, preprocessing throughput grows with worker count as expressed in equation 8.8:

$$R_{\text{preprocessing}} = \frac{N_{\text{workers}}}{T_{\text{transform}}} \quad (8.8)$$

where N_{workers} parallel processing threads each perform transformations requiring $T_{\text{transform}}$ seconds. Training architectures employ multiple workers to ensure preprocessing keeps pace with accelerator consumption rates—a single thread performing image augmentation at 30 ms per batch cannot feed an accelerator that computes a forward pass in 10 ms.

Preprocessed data must then transfer to the accelerator before computation can begin. The overall training throughput is therefore constrained by the slowest of three stages, as equation 8.9 makes explicit:

$$R_{\text{training}} = \min \left(R_{\text{preprocessing}}, \frac{\text{BW}_{\text{GPU transfer}}}{D_{\text{vol}}^{(\text{batch})}}, \frac{R_{\text{compute}}}{O_{\text{batch}}} \right) \quad (8.9)$$

where $D_{\text{vol}}^{(\text{batch})}$ is the input bytes per batch and O_{batch} is the forward-backward work per batch in FLOPs. Every term is therefore expressed in batches per second before the minimum is taken.

Example 8.1: GPT-2 language model data pipeline

Scenario: Quantifying storage access, CPU tokenization, and PCIe host-to-device transfers supplying a 32-V100 GPU cluster training GPT-2.

Stage trace:

1. **Raw text storage:** The WebText corpus contributes about 40 GB of raw text. Sequential reads from NVMe SSD storage reach 7 GB/s. Under document sampling with access factor 0.1, random-access storage bandwidth drops to 0.70 GB/s, giving $T_{\text{read}} = D_{\text{vol}}/\text{BW}_{\text{storage}} \approx 0.1$ ms.
2. **Tokenization:** A BPE tokenizer (50,257 vocabulary) converts text into token IDs. For batch size $B = B = 32$ and sequence length $S = S = 1024$ (32.8K tokens total), a single CPU core processing 500K tokens/s takes $T_{\text{tokenize}} \approx 65.5$ ms per batch.
3. **Batching and padding:** Padding and packing generate an int64 tensor of shape 32 sequences $\times$ 1,024 tokens, producing 262.1 KB per batch.
4. **PCIe transfer:** Transferring the tensor across a PCIe Gen3 x16 bus (15.75 GB/s bandwidth) takes $T_{\text{transfer}} = 0.017$ ms, making bus transfer negligible.

Systems insight: Single-core CPU preparation ($T_{\text{prep}} \approx 65.6$ ms) consumes most of the accelerator step budget ($T_{\text{step}} \approx 84.4$ ms), so a serial pipeline pays both costs on every step and η_{hw} falls to about 56 percent. Push preparation any slower and the accelerator starves outright. Scaling to 8 parallel worker processes reduces tokenization latency to $T_{\text{tokenize}}/N_{\text{worker}} \approx 8.2$ ms, while prefetching overlaps T_{prep} with T_{step} to restore hardware utilization to 95 percent, delivering 379 global samples per second across 32 V100 GPUs.

This min-of-three relationship governs training pipeline design because system throughput equals bottleneck throughput. An accelerator with 312 TFLOP/s of compute capacity delivers zero useful work while waiting for data. Conversely, a perfectly optimized data pipeline provides no benefit if the accelerator is already compute-saturated. Balanced pipeline design aligns preprocessing

capacity, transfer bandwidth, and compute throughput so that no single stage dominates iteration time. Applying this throughput analysis to the GPT-2 lighthouse model reveals where the data pipeline bottleneck lies for language model training.

Data pipeline throughput is not the only flow that can starve the accelerator. Multi-GPU training adds a second stream of traffic, the gradient synchronization required after every step, and once synchronization time exceeds compute time, communication becomes the limiting wall. Section 8.7.2 quantifies that network wall at the point where training crosses the node boundary.

8.4.2.3 System implications

The data pipeline and compute engine form a coupled system whose throughput equals the slower of the two, as equation 8.10 states:

$$R_{\text{system}} = \min(R_{\text{pipeline}}, R_{\text{compute}}) \quad (8.10)$$

This relationship has direct consequences. When $R_{\text{pipeline}} < R_{\text{compute}}$, the accelerator sits idle waiting for data, and accelerator utilization drops proportionally, as equation 8.11 shows:

$$\text{Accelerator Utilization} = \min\left(1, \frac{R_{\text{pipeline}}}{R_{\text{compute}}}\right) \times 100\% \quad (8.11)$$

A ResNet-50 model on modern accelerator hardware can process 1,000 images per second, but if the data pipeline delivers only 200 images per second, accelerator utilization drops to 20 percent—the accelerator is idle 80 percent of the time. Upgrading to faster hardware does not help in this case; an accelerator capable of 2,000 images per second would achieve only 10 percent utilization with the same pipeline. Balanced system design aims to keep the accelerator supplied with data rather than waiting on an upstream stage.

8.4.2.4 Data flows

Training data traverses three memory tiers on its way from disk to accelerator, and the bandwidth gap between these tiers, spanning three orders of magnitude, is the central challenge of data pipeline design. The effective transfer rate through the hierarchy is bounded by its slowest link, as equation 8.12 shows:

$$R_{\text{memory}} = \min(\text{BW}_{\text{storage}}, \text{BW}_{\text{system}}, \text{BW}_{\text{accelerator}}) \quad (8.12)$$

Local NVMe storage provides 7 GB/s, system memory delivers 50 GB/s, and accelerator HBM achieves 900 GB/s or higher. Each tier is roughly an order of magnitude faster than the one below it, which means data that flows freely within accelerator memory creates a severe bottleneck when it must be fetched from disk. This cascading bandwidth hierarchy explains why the iteration time of a well-pipelined system is governed by the *maximum* of its component latencies rather than their *sum*, as equation 8.13 shows:

$$T_{\text{iteration}} = \max(T_{\text{fetch}}, T_{\text{process}}, T_{\text{transfer}}) \quad (8.13)$$


When pipeline stages overlap correctly (fetching the next batch from storage while preprocessing the current one and transferring the previous one to the accelerator), the iteration time equals the duration of the slowest stage rather than the sum of all stages. This overlap is exactly what prefetching achieves, turning a serial bottleneck into a parallel pipeline where each tier operates concurrently on different batches.

8.4.2.5 Practical architectures

An NVMe device rated at 7 GB/s would sustain 3.50 GB/s if a workload achieved half of peak. Small random reads and shuffling can reduce throughput further; the penalty depends on access size, queue depth, file layout, and caching.

To keep accelerators fed despite this bandwidth reduction, pipeline architectures maintain multiple data buffers simultaneously—prefetch buffers loading future batches, processing buffers holding data under transformation, and transfer buffers staging data for accelerator consumption. The total host memory required scales with the per-batch memory footprint M_{batch} according to equation 8.14:

$$M_{\text{required}} = (N_{\text{prefetch}} + N_{\text{processing}} + N_{\text{transfer}}) \times M_{\text{batch}} \quad (8.14)$$



Bandwidth steps up the storage to DRAM to HBM hierarchy.

To avoid starving the accelerator, average preprocessing service time per batch must satisfy equation 8.15.

$$T_{\text{preprocessing}} \leq T_{\text{compute}} \quad (8.15)$$

If preprocessing takes 40 ms per batch on one worker while the accelerator consumes one every 10 ms, ideal scaling requires at least four workers. Actual requirements depend on storage, CPU contention, and transformation cost; section 8.6.2 examines prefetching and parallel workers.

8.4.3 Forward pass

Prepared batches enter the training loop through the forward pass, where input data propagates through the model to generate predictions. The conceptual flow follows the layer-by-layer transformation $\mathbf{A}^{(\ell)} = f(\mathbf{A}^{(\ell-1)}\mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)})$ established earlier, but the system-level implementation must schedule kernels, move activations, and preserve enough state for backpropagation.

8.4.3.1 Compute operations

The forward pass orchestrates the computational patterns introduced in section 8.3.1.2, optimizing them for specific neural network operations. Building on the matrix multiplication foundations, the system must efficiently execute the $N \times M \times B$ multiply-accumulate operations required for each layer, where typical layers with dimensions of 512×1024 processing batches of 64 samples execute about 33.6 million MACs, or about 67.1 MFLOP under the two-FLOPs-per-MAC convention.

Systems Perspective 8.6: Wave quantization and tail effects

A common mistake in ML systems is ignoring the fixed execution granularity of GPU kernels. GPU execution is quantized into “waves” of work, but the relevant unit is the kernel’s thread or tile mapping, not necessarily one sample per thread.

The wave effect: An NVIDIA GPU executes work in warps of 32 threads. In the illustrative special case where one independent work item maps to one thread, 32 items fill one warp while 33 items require a second, mostly empty warp. Neural-network kernels usually map each sample to many threads, so batch size alone does not determine this utilization.

Tail effects at scale: On a large GPU like the H100 with 132 Streaming Multiprocessors (SMs), where each SM schedules groups of warps, the hardware can process thousands of threads in one “wave.” If the total workload is just slightly over a wave boundary (e.g., 1.01 waves), the hardware must wait for a nearly empty wave to finish before the next task begins.

Table 8.6 quantifies that one-item-per-thread illustration. It is not a general batch-size performance model.

Table 8.6: Illustrative Warp-Tail Arithmetic: Under a one-work-item-per-thread mapping, counts just above a 32-thread boundary launch a partially filled warp, so 33 items use two warps while 64 fill the same two warps. Actual training kernels map tensors to threads and tiles in implementation-specific ways, often assigning many threads to each sample, so these values cannot be inferred from batch size alone. Profiling must identify the kernel’s real mapping and wave boundaries.

Independent Work Items	Warps Needed	Lane Utilization	Relative Time
32	1	100%	1×
33	2	51.6%	~2×
64	2	100%	1×
65	3	67.7%	~1.5×

Engineering rule: Align tensor and tile dimensions with the kernel’s hardware mapping, then benchmark candidate batch sizes. A profiler can reveal partially filled warps or waves; batch size by itself cannot.

Understanding these tail effects is the difference between a practitioner who tunes by trial-and-error and an engineer who designs for the hardware.

Modern neural architectures extend beyond these basic matrix operations to include specialized computational patterns. Convolutional networks, for instance, perform systematic kernel operations across input tensors. Consider a typical input tensor of dimensions $64 \times 224 \times 224 \times 3$ (batch size $\times$ height $\times$ width $\times$ channels) processed by 7×7 kernels. Each position requires 147 multiply-accumulate operations, and with 64 filters operating across 218×218 spatial dimensions, the computational demands become substantial.

Transformer architectures introduce attention mechanisms (see chapter 6), which compute similarity scores between sequences. These operations combine matrix multiplications with softmax normalization, requiring efficient broadcasting and reduction operations across varying sequence lengths. The computational pattern here differs significantly from convolutions, demanding flexible execution strategies from hardware accelerators.

Throughout these networks, elementwise operations play a supporting role. Standalone implementations of activation functions such as ReLU and sigmoid transform values independently and are often memory-bandwidth-bound rather than compute-bound (see section 8.3.4). Batch normalization presents similar challenges, computing statistics and normalizing values across batch dimensions while creating synchronization points in the computation pipeline.

Modern hardware accelerators, particularly GPUs, optimize these diverse computations through massive parallelization. Achieving peak performance requires careful attention to hardware architecture. GPUs process data in fixed-size groups of threads called warps¹⁹ on NVIDIA architectures or wavefronts on AMD architectures. Efficient matrix dimensions depend on the kernel's thread-block and tile mapping, including any alignment requirements imposed by Tensor Core instructions.

Libraries like cuDNN (Chetlur et al. 2014) address these challenges by providing optimized implementations for each operation type. These systems dynamically select algorithms based on input dimensions, hardware capabilities, and memory constraints. The selection process balances computational efficiency with memory usage, often requiring empirical measurement to determine optimal configurations for specific hardware setups (NVIDIA Corporation 2023). These hardware utilization patterns reinforce the batch-size–utilization relationship established in section 8.3.2.1: the tension between larger batch sizes (better utilization) and memory constraints (forcing smaller batches) permeates all levels of training system design.

8.4.3.2 Memory management

Memory management binds during the forward pass, when intermediate activations must be stored for subsequent backward propagation. Before examining how frameworks manage forward-pass memory, it is useful to estimate the total VRAM required for training. A concrete 7-billion-parameter-on-24 GB case shows how weights, gradients, optimizer state, and activations combine.



Napkin Math 8.4: Estimating VRAM requirements

Problem: Will a 7B-parameter model fit on a 24 GB GPU for training?

Given: 7B parameters, FP16 weights and gradients, FP32 master weights and Adam states, and 24 GB GPU memory.

Math:

1. **Weights (FP16):** $7\text{B} \times 2 \text{ bytes} = 14 \text{ GB}$.
2. **Gradients (FP16):** Same size as weights = 14 GB.
3. **FP32 training state:** Master weights require 28 GB; Adam momentum and variance require another 56 GB.
4. **Subtotal (before activations):** $14 \text{ GB} + 14 \text{ GB} + 28 \text{ GB} + 56 \text{ GB} = 112 \text{ GB}$, already exceeding a 24 GB GPU.
5. **Activations:** A simplified transformer estimate is $\text{Batch} \times \text{SeqLen} \times \text{Hidden} \times \text{Layers} \times \text{Bytes} \times \text{an activation factor for retained attention, multilayer perceptron, and normalization intermediates}$. With $\text{Batch} = 1$, $\text{Seq} = 2048$, $\text{Hidden} = 4096$, 32 layers, and an activation factor of about $57\times$, retained activations add 30.6 GB.

¹⁹ | **Warp:** From textile weaving, where a “warp” is the set of threads held taut on a loom while a shuttle weaves across them. NVIDIA adopted the term because GPU threads within a warp execute the same instruction in lockstep, analogous to parallel threads moving together on the loom. An NVIDIA warp contains 32 threads.

Systems insight: Full-parameter training requires partitioning or offloading state through systems such as fully sharded data parallel (FSDP) or ZeRO. Quantized base weights reduce parameter-efficient fine-tuning memory but do not reproduce this full-parameter mixed-precision recipe.

Only activation memory scales approximately linearly with batch size (equation 8.5); model and optimizer state are fixed. The practical complexity lies in how these costs interact across layers.

Consider a representative large model like ResNet-50 (a widely-used image classification architecture) processing images at 224×224 resolution with a batch size of 32. The initial convolutional layer produces activation maps of dimension $112 \times 112 \times 64$; per image at single-precision (4 bytes), this requires approximately 3.2 MB. As the network progresses through 50 layers, the cumulative memory demands grow substantially: the complete forward pass activations total approximately 8 GB, the backward pass adds another 4 GB of activation working set (empirical total, not stored parameter gradients), and model parameters consume 102.4 MB. This 12.1 GB total represents about 15.1 percent of an A100 GPU's 80 GB memory capacity for a single batch.

The memory scaling patterns reveal critical hardware utilization trade-offs. Doubling the batch size to 64 increases forward activation memory to 16 GB and the backward activation working set to 8 GB, totaling 24.1 GB and reducing the memory headroom available for deeper models, larger inputs, and optimizer state. Training larger models at the scale of GPT-3 (175B parameters, representing current large language models) requires approximately 700 GB just for parameters in FP32 (350 GB in FP16), necessitating distributed memory strategies across multiple high-memory nodes.

GPUs typically provide 40 GB–80 GB of memory in high-end training configurations, which must accommodate activations, model parameters, gradients, and optimization states. Two techniques address this constraint directly: activation checkpointing trades recomputation for reduced activation storage, and mixed-precision training halves memory per value by using FP16 instead of FP32. Both are examined in detail in section 8.6; here, the key insight is that memory capacity, not compute throughput, often determines the maximum feasible batch size and model depth. Practitioners frequently start with large batch sizes during initial development on smaller networks, then adjust downward when scaling to deeper architectures or memory-constrained hardware.

The backward pass reverses this flow, computing gradients at approximately twice the forward pass cost (section 8.3.3). The per-layer memory costs accumulate rapidly across the full network: deeper in ResNet-50, mid-network convolutional layers use 256 filters rather than the initial 64, but smaller spatial maps offset some activation memory while convolutional work depends on kernel size and both input and output channels. Across 50 layers, the illustrative backward-pass working set reaches approximately 3.2 GB before accounting for optimizer state and parameter updates. Dependencies order the layerwise gradient computations, although runtimes may overlap independent kernels and gradient communication. Peak memory reflects retained forward activations, gradients, temporary workspaces, and the framework's allocation schedule rather than the width of a single layer alone.

8.4.4 Parameter updates and optimizers

Once the backward pass computes gradients, the system must allocate and manage memory for both parameters and gradients, then perform the update computations. The choice of optimizer determines the mathematical update rule and the system resources required for training. Listing 8.3 demonstrates the backward/update portion of the parameter update cycle in PyTorch: after the forward pass computes predictions and the loss function quantifies error, `loss.backward()` populates gradient tensors and `optimizer.step()` applies the update rule to all parameters based on the configured optimizer (Adam, SGD, etc.).

Listing 8.3: Parameter Update: Computes gradients and applies an optimizer step to model parameters. Repeating this cycle seeks to minimize the training objective, although improvement is not guaranteed on every step or epoch.

```
loss.backward() # Compute gradients
optimizer.step() # Update parameters
```

These operations initiate a sequence of memory accesses and computations. The system must load parameters from memory, compute updates using the stored gradients, and write the modified parameters back to memory. Different optimizers vary in their memory requirements and computational patterns, directly affecting system performance and resource utilization.

8.4.4.1 Optimizer memory in the training loop

The optimizer memory hierarchy established in table 8.4 manifests concretely during each training iteration. Each parameter update involves reading current values, accessing gradients, computing the update rule, and writing modified parameters back to memory. For Adam, this includes updating and accessing the momentum and variance buffers, creating substantial memory traffic for large models.

At billion-parameter scale, optimizer state dominates the memory budget. As quantified in the GPT-2 worked example (section 8.3.2.3), a 1.5B model requires 12 GB for Adam optimizer state alone in FP32, in addition to parameters and gradients, before accounting for activations. This challenge has motivated memory-efficient optimizer variants. Adafactor factorizes second-moment state (Shazeer and Stern 2018), 8-bit optimizers quantize optimizer statistics (Dettmers et al. 2022), and GaLore computes updates in a low-rank space. Compare the memory bars in figure 8.6 to see how GaLore attacks this constraint: by computing updates in a compressed space (Zhao et al. 2024), the technique reduces the memory footprint dominated by optimizer states to a fraction of its original size, enabling training of larger models on fixed hardware.




Figure 8.6: Illustrative Memory Footprint Breakdown: Memory usage of a 7-billion-parameter LLaMA model across four memory-efficient optimizer configurations (BF16 Adam, Adafactor, 8-bit Adam, and 8-bit GaLore), decomposed into weights, activations, optimizer state, weight gradients, and other components. The dashed red line marks the RTX 4090 24 GB memory limit. Standard FP32 Adam (omitted from the bars; it requires well above 70 GB for this model) does not fit on a single 24 GB GPU; the bars compare relative memory reductions, but even the smallest illustrated total remains above the 24 GB line, with 8-bit GaLore shrinking optimizer state most aggressively.

The bars make the ranking concrete: BF16 Adam pushes farthest past the single-GPU budget line, and Adafactor, 8-bit Adam, and 8-bit GaLore progressively reduce the overage. Even the smallest illustrated total remains above 24 GB, so another memory-saving technique is still required.

8.4.4.2 Batch size and parameter updates

Scaling batch size without retuning the learning-rate schedule can degrade convergence. Compare the illustrative loss trajectories in figure 8.7, observing how the fixed learning rate slows convergence relative to the scaled run in this scenario.

For a fixed dataset, doubling the global batch size halves the number of updates per epoch. The **linear scaling rule** ($\eta_{\text{new}} = k \times \eta_{\text{base}}$) is a useful empirical recipe over some batch-size ranges, often combined with learning-rate warmup, but it does not guarantee equivalent convergence or generalization. Figure 8.7 uses normalized synthetic curves to illustrate the intended effect.

Beyond the convergence effects, batch size interacts with distributed training strategies. A larger global batch reduces the number of optimizer steps and gradient synchronizations per epoch, while




Figure 8.7: Illustrative Linear-Scaling Scenario: Synthetic training-loss curves (arbitrary units) compare a batch-32 baseline with batch 256 under a fixed or linearly scaled learning rate. The curves illustrate why increasing batch size may require retuning; they are not measured traces or a guarantee that linear scaling preserves convergence.

the payload of each dense synchronization is primarily determined by the parameter-gradient size rather than the batch size. In distributed settings, local and global batch sizes constrain the degree of data parallelism and the frequency of parameter updates. Gradient accumulation (section 8.6.5) decouples the effective batch size from the number of samples held for one micro-batch, though it does not remove the need to tune the effective batch size.

 Napkin Math 8.5: The utility bill

Problem: Under the following simplified assumptions, is it cheaper to rent or purchase a 1,024-H100 cluster for one Llama 2 70-billion-parameter training run?

Math:

1. Workload: Llama 2 70B model (70B parameters, 2T tokens).
2. **Compute required:** $6 \times 70 \times 10^9 \times 2 \times 10^{12} \approx 8.4 \times 10^{23}$ FLOPs. The leading factor of 6 follows from the chapter’s accounting: about 2 FLOPs per parameter per token in the forward pass (two FLOPs per multiply-accumulate), tripled once the backward pass, which costs roughly twice the forward pass, is included.
3. Hardware: NVIDIA H100 (Peak: 989 TFLOP/s FP16). Assumed Utilization: 50 percent (494.5 TFLOP/s).
4. Time: $8.4 \times 10^{23} / (494.5 \times 10^{12}) \approx 1.70 \times 10^9$ seconds $\approx$ 53.8 years (on one GPU).
5. **Cluster:** On 1,024 GPUs $\rightarrow$ 19.2 days.

The economics:

- **Rental (\$3/hr):** 1,024 GPUs $\times$ 24 hrs $\times$ 19.2 days $\times$ \$3/hr $\approx$ \$1.42M.
- **Purchase (\$30,000 per GPU):** 1,024 GPUs $\times$ \$30,000 = \$30.7M.

Systems insight: In this simplified utilization-and-price scenario, purchase cost equals about 21.7 rental runs. A real ownership comparison must also include cluster networking, host systems, facilities, power, staffing, depreciation, financing, and the utilization achieved between runs.

Every batch-size and precision decision so far has aimed at wall-clock time; at scale, that time is denominated in dollars. The compute cost itself becomes a binding constraint that shapes every training decision, from hardware selection to cluster sizing. The calculation here turns that cost into a rental-versus-purchase decision for a realistic training run.

This section established the structural *what* of training systems, and the mathematical foundations in section 8.3 quantified the FLOPs, memory, and bandwidth each stage demands. Yet understanding what must happen does not reveal *where* the system currently underperforms. A training pipeline is only as fast as its slowest stage: if data loading takes 50 ms and computation takes 100 ms, optimizing computation by 20 percent saves 20 ms, but if the bottleneck were data loading, those same engineering hours would save nothing. Before reaching for optimization techniques, diagnostic tools must identify which constraint actually limits performance.

8.5 Identifying Bottlenecks

Blueprint knowledge is not diagnosis. Knowing that attention operations consume 50 percent of FLOPs and data loading takes 25 percent of wall-clock time does not reveal which constraint to attack first; that depends on which resource is actually saturated during execution.

The diagnostic methodology that transforms blueprint knowledge into actionable optimization decisions begins with a meaningful measure of training efficiency. Raw accelerator utilization percentages can be misleading because an accelerator may remain active while executing recomputation, padding, or other work outside the model-FLOP accounting. Model FLOPs utilization (MFU)²⁰ supplies that comparison.

²⁰ | **Model FLOPs utilization (MFU):** The PaLM paper (Chowdhery et al. 2022) defined MFU as observed throughput divided by the theoretical maximum throughput at peak FLOPs. Its 540-billion-parameter run reported 46.2 percent MFU on 6,144 TPU v4 chips. The complement is uncredited by the model-FLOP numerator, but MFU alone does not identify how much came from memory traffic, communication, pipeline bubbles, recomputation, or other work.



Definition 8.4: Model FLOPs utilization (MFU)

Model FLOPs utilization (MFU) is the efficiency metric $MFU = O_{\text{model}} / (R_{\text{peak}} \cdot T_{\text{step}})$, where O_{model} is the useful per-step model FLOP count (forward plus backward, excluding rematerialization) and T_{step} is the measured wall-clock time per training step, expressing what fraction of peak hardware throughput is doing useful model computation.

1. **Significance:** MFU makes the effective-throughput term in the iron law concrete. For a 7B-parameter transformer processing 1,024 tokens per step on an A100 (312 TFLOP/s FP16/BF16 Tensor Core peak), a 1.2-second step gives model FLOPs divided by the available peak-rate FLOP budget of approximately 0.11 (11.5 percent). The remaining 88.5 percent is not credited as model FLOPs and can reflect memory traffic, communication, bubbles, non-model operations, or idle time. Whether a value is good depends on the model, hardware, precision, and parallelization strategy; diagnosis requires a profile rather than a universal threshold.
2. **Distinction:** Hardware utilization reports whether device engines are active, whereas MFU credits the analytical model-FLOP estimate used in its numerator. It can therefore exclude recomputation and padding even when those operations keep the device busy. Comparisons across accelerators remain sensitive to the selected peak-throughput specification and FLOP-counting convention.
3. **Common pitfall:** A frequent misconception is that 100 percent MFU is a realistic sustained target. Memory traffic, communication, and non-matrix work keep end-to-end runs below peak, but no universal 55–65 percent ceiling exists. Reported values depend on the model, precision, batch shape, hardware, and MFU definition, and can improve with FlashAttention and careful tuning.

Training bottlenecks can be investigated through three categories aligned with the D-A-M taxonomy (Data, Algorithm, Machine; Appendix A provides the full diagnostic framework, troubleshooting matrix, and D-A-M Scorecard). Table 8.7 connects each axis to a common training bottleneck, its observable symptoms, and relevant optimization techniques. These are diagnostic starting points, not exclusive mappings.

Table 8.7: D-A-M Taxonomy Applied to Training Bottlenecks: Each row associates a D-A-M axis with one common bottleneck and its characteristic symptoms. Profiling reveals the workload’s limiting constraint and guides practitioners to an appropriate optimization technique.

D-A-M Axis	Bottleneck	Symptoms	Primary Solutions
Algorithm	Compute-bound	High arithmetic-unit activity; additional bandwidth does not improve throughput; arithmetic throughput is limiting	Reduced precision, more efficient algorithms, faster compute hardware
Machine	Memory-bound	High memory-bandwidth use relative to arithmetic throughput; kernels stall on data movement	Operator fusion, memory-efficient attention, reduced precision formats
Data	Data-bound	Periodic accelerator utilization drops to near-zero; CPU fully busy during gaps; pipeline cannot feed GPU fast enough	Prefetching, pipeline overlap, faster storage, DataLoader parallelism

The data-bound category is a commonly misdiagnosed bottleneck in practice.

 Example 8.2: The GIL-locked GPU

Scenario: An engineering team builds a PyTorch data-loading pipeline using standard Python `threading` for CPU-bound image preprocessing that holds the global interpreter lock (GIL).

Diagnosis: The GIL serializes these Python-level preprocessing threads, making host preparation time ($T_{\text{prep}} \approx 250$ ms) vastly exceed GPU execution time ($T_{\text{step}} \approx 12$ ms), collapsing hardware utilization to $\eta_{\text{hw}} \approx 4.6\%$.

Systems lesson: Python CPU work that holds the GIL can serialize threads and starve GPUs; native operations may release it. Process workers (`num_workers > 0`) avoid interpreter serialization, while NVIDIA DALI can move eligible work off the host.

Profiling tools reveal which bottleneck dominates a given workload. Figure 8.8 captures the data-bound pathology from the callout: white gaps between repeated compute blocks mark intervals with no GPU kernels visible in the device timeline.




Figure 8.8: Data-Bound Profiler Trace: In this timeline trace, isolated GPU compute blocks are separated by prominent white intervals. These idle gaps mark periods where accelerator streams stall waiting for CPU-side batch preparation and host-to-device transfers, providing a characteristic signature of an input-bound pipeline.

Four tools integrated into machine learning frameworks provide detailed bottleneck analysis. Table 8.8 separates framework-level timeline tools from GPU-level execution tools, because each exposes a different failure signature.

Table 8.8: Profiling Tools for Training Bottlenecks: Framework profilers expose operation timelines and input-pipeline stalls, while NVIDIA Nsight tools expose system-level GPU execution and kernel-level memory behavior.

Tool	Scope	Best signal
PyTorch Profiler (<code>torch.profiler</code>)	Framework-level operation trace	Time spent in each operation, memory allocation patterns, and GPU kernel execution
TensorFlow Profiler	Framework-level training timeline	Input pipeline bottlenecks and device placement
NVIDIA Nsight Systems	System-level GPU trace	Kernel execution, memory transfers, and synchronization points
NVIDIA Nsight Compute	Kernel-level GPU analysis	Arithmetic intensity, memory throughput, and occupancy

The profiling workflow follows a systematic pattern: run a representative training iteration with profiling enabled, examine the timeline for gaps (data-bound), check memory bandwidth utilization (memory-bound vs. compute bound), and identify the dominant bottleneck before selecting an optimization technique.

In practice, the characteristic signatures from table 8.7 are directly visible in profiler traces: accelerator utilization levels, memory bandwidth saturation, and CPU vs. GPU activity ratios each point to a specific bottleneck class. These signatures map to specific optimization techniques: prefetching for data bottlenecks, mixed precision and operator fusion for memory bottlenecks, and algorithmic improvements or hardware upgrades for compute bottlenecks. With a diagnostic framework in hand, the next step is to examine each optimization technique in detail: what it does, which iron law term it targets, and when profiling results indicate it should be applied.

8.6 Pipeline Optimizations

Profiling reveals *where* the training system underperforms; the D-A-M taxonomy (Appendix A) classifies *what kind* of bottleneck limits throughput. The remaining question is *how* to close the gap. Four optimization techniques, each targeting a specific bottleneck category, together with a systematic framework for composing them, provide the answer.

Even well-designed pipeline architectures rarely reach hardware peak throughput without targeted optimization. To illustrate the scale of the gap, applying a model-FLOP utilization of 30 percent to 50 percent to an A100 peak of 312 TFLOP/s gives 93.6 TFLOP/s–156 TFLOP/s of model-FLOP throughput. These are derived scenario values, not measurements reported for one A100 workload. Published large-model systems report similar gaps between peak and realized throughput (Narayanan et al. 2021; Chowdhery et al. 2022). Systems such as SuperNeurons show one version of this problem: memory-management bottlenecks can prevent larger networks from training efficiently unless the runtime actively optimizes activation storage and transfer (L. Wang et al. 2018). Table 8.9 extends the D-A-M-based bottleneck classification from table 8.7 by mapping each bottleneck to the specific optimization technique that addresses it.

Table 8.9: Optimization Technique Roadmap: Each primary bottleneck category has targeted solutions that address specific performance constraints, matching techniques to profiling results for systematic optimization.

Bottleneck	Primary Solution(s)
Data Movement Latency	Prefetching & Pipeline Overlapping
Compute Throughput	Mixed-Precision Training
Memory Capacity	Gradient Accumulation & Activation Checkpointing
Memory Bandwidth (Attn.)	FlashAttention (IO-aware tiling)

These bottlenecks manifest differently across system scales (a 100 GB model faces different constraints than a 1 GB model), but identification and mitigation follow consistent principles. Data movement latency emerges when training batches cannot flow from storage through preprocessing to compute units fast enough to keep accelerators in use. Computational throughput limitations occur when mathematical operations execute below hardware peak performance due to suboptimal precision choices or kernel inefficiencies. Memory capacity constraints restrict both the model sizes and batch sizes a system can process, directly limiting model complexity and training efficiency.

These bottlenecks interact, illustrating the conservation of complexity thesis from Part I: relieving one constraint can expose another. When data loading becomes a bottleneck, GPUs sit idle waiting for batches. When memory capacity is constrained, smaller batches may reduce GPU efficiency. For a hypothetical GPT-2 profile, suppose attention occupies 50 percent of time, data loading 25 percent, and other compute-bound operations 25 percent. Such a profile motivates evaluating memory-efficient attention, prefetching, and reduced precision, then measuring again to see which change matters. The optimization challenge is to identify the current bottleneck and select techniques that address it without creating a worse constraint elsewhere.

8.6.1 Systematic optimization framework

The profile example shows why optimization must start from evidence rather than preference. Effective optimization follows a systematic methodology that applies regardless of system scale or model architecture: profile to identify bottlenecks, select appropriate techniques for the identified constraints, and compose solutions that address multiple bottlenecks simultaneously without creating conflicts.

The profiling phase employs tools like PyTorch Profiler, TensorFlow Profiler, or NVIDIA Nsight Systems to reveal where time is spent during training iterations. These are the same profiling approaches introduced in the overview, now applied systematically to quantify which bottleneck dominates. A profile might show 40 percent of time in data loading, 35 percent in computation, and 25 percent in memory operations, indicating data loading as the primary target for optimization.

The selection phase matches optimization techniques to identified bottlenecks. Each technique targets specific constraints. Prefetching addresses data movement latency, mixed-precision training tackles both computational throughput and memory constraints, and gradient accumulation manages memory limitations. Selection requires understanding the bottleneck type alongside the characteristics of the hardware, model architecture, and training configuration that influence technique effectiveness.

The composition phase combines multiple techniques to achieve cumulative benefits. Prefetching and mixed-precision training complement each other (one addresses data loading, the other computation and memory), allowing simultaneous application. However, some combinations create conflicts: aggressive prefetching increases memory pressure, potentially conflicting with memory-constrained configurations. Successful composition requires understanding technique interactions and dependencies.

This systematic framework (profile, select, compose) applies to the four core optimization techniques covered next. Prefetching targets data movement latency. Mixed-precision training addresses both throughput and memory constraints. FlashAttention reduces memory traffic for attention. Gradient accumulation manages batch-memory limits by serializing micro-batches, while checkpointing trades recomputation for lower activation storage. The profile should determine the order in which these techniques are evaluated, and each requires a cost-benefit analysis that includes implementation and debugging effort.

Figure 8.9 provides a decision tree that operationalizes this systematic framework. The branches lead from profiling results through bottleneck identification to technique selection, ensuring optimization effort targets the actual constraint rather than perceived issues.

The flowchart embodies a critical insight: optimization is iterative. After applying a technique, re-profiling often reveals that a different bottleneck has become dominant. A data-bound system that implements prefetching may become memory bound, requiring the next technique in the decision tree. This iterative refinement continues until profiling shows balanced resource utilization or acceptable training throughput.

8.6.2 Data prefetching and overlapping

Prefetching and overlapping techniques illustrate the systematic framework in action, targeting data movement latency bottlenecks by coordinating data transfer with computation. This optimization proves most effective when profiling reveals that computational units remain idle while waiting for data transfers to complete.

Training machine learning models involves significant data movement between storage, memory, and computational units. The data pipeline consists of sequential transfers: from disk storage




Figure 8.9: Training Optimization Decision Flowchart: A heuristic path from profiling signals to candidate optimizations. Begin with accelerator utilization, then distinguish memory pressure from compute saturation before selecting an intervention. The 70 percent utilization and 90 percent memory thresholds are illustrative triage values rather than universal definitions of data-, memory-, or compute-bound execution.

to CPU memory, CPU memory to GPU memory, and through the GPU processing units. In ML training, a “Read” operation is rarely a simple disk fetch: for image workloads it encompasses JPEG decoding, random crops, and color jitter applied on CPU cores before the tensor is ready to transfer; for language workloads it encompasses tokenization, subword encoding, and sequence padding. Figure 8.10 exposes the inefficiency of sequential data transfer: the GPU remains idle during file operations (Open 1, Open 2), and training steps cannot begin until these read and preprocessing operations complete, leaving expensive compute resources underutilized for significant portions of each epoch.

Prefetching addresses these inefficiencies by loading data into memory before its scheduled computation time. During the processing of the current batch, framework data pipelines such as `tf.data` load and prepare subsequent batches, maintaining a consistent supply of ready data (Murray et al. 2021).




Figure 8.10: Illustrative Sequential Data Fetching: Under the assumed stage durations, file-open, read, and train operations execute serially while the GPU remains idle during file operations. The schedule spans approximately 95 min and provides a schematic baseline for the overlap comparison.

Overlapping extends prefetching by coordinating multiple pipeline stages to execute concurrently. The system processes the current batch while simultaneously preparing future batches through data loading and preprocessing operations. Compare the illustrative schedules in figure 8.10 and figure 8.11. Under their assumed stage durations, overlap reduces completion time from approximately 95 min to 65 min, a 31.6 percent reduction. Actual gains depend on stage balance and available concurrency.




Figure 8.11: Illustrative Overlapped Data Prefetching: File-open and read work overlaps accelerator training while preserving the same individual stage durations as the serial schedule. The schematic schedule spans approximately 65 min, compared with 95 min for the serial schedule.

8.6.2.1 Prefetching mechanics

Prefetching only pays off when the profile shows data movement on the critical path. Training data still passes through retrieval, transformation, and model execution; the optimization changes when those stages run. An unoptimized pipeline serializes them, leaving the GPU idle during data fetching and preprocessing. Data loaders avoid that stall by preparing the next batch in separate threads or processes while the current batch trains.

Overlapping extends the same idea across the full pipeline. As the GPU processes one batch, preprocessing begins on the next batch, while data fetching starts for the subsequent batch. The goal is constant activity at each stage so the slowest stage no longer leaves the others waiting.

Machine learning frameworks (introduced in chapter 7) expose the control variables for this trade-off. Listing 8.4 demonstrates PyTorch’s DataLoader configuration, where `num_workers = 4` sets preprocessing parallelism and `prefetch_factor = 2` maintains a buffer of 8 batches ready for accelerator consumption.

Listing 8.4: DataLoader Prefetch Configuration: Four workers with a prefetch factor of two maintain up to eight batches in the host-side prefetch window.

```
loader = DataLoader(
    dataset, batch_size=32, num_workers=4, prefetch_factor=2
)
```

The parameters `num_workers` and `prefetch_factor` are not generic performance toggles. They determine CPU parallelism and buffer depth, which means they shift pressure from accelerator idle time to host memory and CPU scheduling.

Buffer management is the main trade-off. A buffer that is too small causes the GPU to wait for data preparation, reintroducing the idle time prefetching was meant to remove. An overly large buffer consumes memory that could otherwise store model parameters or larger batches. The right configuration coordinates CPU preparation, storage I/O, and GPU computation so each resource has work ready at the moment it can use it. These techniques yield the greatest benefit when storage access is slow, preprocessing is complex, or datasets are large.

8.6.2.2 Practical considerations

Prefetch buffers and overlap depth are tunable, so the same machinery adapts to whichever stage binds, whether slow storage, limited network bandwidth, or computational throughput. Prefetching

and overlapping deliver the greatest gains when preprocessing lies on the critical path but can run concurrently with model computation. In an illustrative image pipeline, random cropping (10 ms), color jittering (15 ms), and normalization (5 ms) total 30 ms per batch. Overlap can hide up to the concurrent accelerator-computation time; any excess remains on the critical path. NLP workloads similarly benefit when tokenization and subword processing would otherwise block the training loop.

The primary trade-off is memory: prefetch buffers consume host memory, or device memory when a separate device-side queue is used, in proportion to buffer depth and batch size. With batch size 256 high-resolution images (1024×1024 pixels), one buffered batch requires approximately 3.2 GB. With `num_workers = 4` and `prefetch_factor = 2`, the 8-batch prefetch window can hold about 25.8 GB. Tuning these parameters requires empirical testing because excessive workers contend for CPU, storage, and memory resources, while insufficient buffering reintroduces data stalls. Start with a modest worker count, increase it while measuring throughput and memory, and stop when additional workers no longer reduce accelerator idle time. When the input pipeline already exceeds compute demand, deeper prefetching adds complexity without improving throughput.

8.6.3 Mixed-precision training

While prefetching optimizes data movement, mixed-precision training addresses both computational throughput limitations and memory capacity constraints. It uses a lower-precision format where appropriate while retaining higher precision for selected accumulations or updates. Section D.4 compares FP32, FP16, BF16, FP8, and INT8 precision-range trade-offs. A training recipe commonly combines FP32 with either FP16 or BF16; it does not ordinarily use FP16 and BF16 interchangeably in the same path. The resulting speed, memory, and accuracy effects depend on the workload and hardware (Micikevicius et al. 2017; Wang and Kanwar 2019; Kalamkar et al. 2019).

A neural network’s stored FP32 weights require 4 bytes per parameter, while FP16 and BF16 weights use 2 bytes. For a model with 10^9 parameters, the weight tensor shrinks from 4 GB to 2 GB. Total training memory falls by less when the recipe retains FP32 master weights, optimizer states, or selected activations.

The numerical differences between these formats shape their use cases. Table 8.10 shows that BF16’s 8-bit exponent gives it approximately the same normal exponent range as FP32, while FP16’s 5-bit exponent gives a minimum positive normal value of about 6.1×10^{-5} and a maximum finite value of 65,504. FP16 subnormals extend toward 6×10^{-8} , although their treatment depends on the operation and hardware mode. FP32 provides about seven decimal digits of precision, FP16 about three, and BF16 about two to three. BF16 therefore trades mantissa precision for the exponent range that often simplifies deep-learning training.

Table 8.10: Precision Format Comparison: The choice between FP16 and BF16 depends on whether dynamic range (BF16’s strength) or precision (FP16’s advantage) matters more for the specific workload. Minimum normal values shown are the practical thresholds for training, as subnormal values may flush to zero on many GPUs. On A100, FP16 and BF16 Tensor Cores reach $16\times$ the FP32 CUDA-core peak.

Property	FP32	FP16	BF16
Exponent bits	8	5	8
Mantissa bits	23	10	7
Min normal value	10^{-38}	6.1×10^{-5}	10^{-38}
A100 peak throughput ratio	$1\times$	$16\times$	$16\times$

The choice between formats depends on model characteristics and hardware support. BF16’s wider exponent range can reduce overflow and underflow concerns, while FP16 provides more mantissa bits but often requires loss scaling. Framework autocast policies and fused kernels choose higher-precision accumulation for numerically sensitive operations when needed; loss reduction, softmax, normalization, and optimizer updates are common candidates, but their exact execution dtype is implementation-dependent.

Figure 8.12 traces a conventional FP16 mixed-precision recipe. FP32 master weights are cast for lower-precision computation, the loss is scaled before backpropagation, and gradients are unscaled

before the FP32 update. Modern framework implementations may store or accumulate particular tensors differently, and BF16 recipes commonly omit loss scaling because of their wider exponent range.




Figure 8.12: Conventional FP16 Mixed-Precision Training: A schematic six-step recipe casts FP32 master weights for lower-precision computation, scales the loss, computes scaled gradients, then unscales them for an FP32 update. Framework implementations may vary, and BF16 usually does not require this loss-scaling path.

Modern hardware architectures are specifically designed to accelerate reduced precision computations. NVIDIA Tensor Cores were introduced for FP16 mixed-precision operations (NVIDIA 2017), and later A100-class Tensor Cores added BF16 support (NVIDIA Corporation 2020a). Google’s TPUs natively support BF16, as this format was specifically designed for machine learning workloads (Wang and Kanwar 2019). These architectural optimizations typically enable substantially higher computational throughput for reduced precision operations compared to FP32, making mixed-precision training particularly efficient on modern hardware.

Mixed-precision training turns those hardware capabilities into a split numerical contract. Eligible matrix multiplications and convolutions use FP16 or BF16 inputs on Tensor Core paths, while selected reductions, accumulations, and parameter updates retain higher precision. The performance gain comes from both arithmetic and data movement. Lower-precision kernels can execute at higher throughput, and smaller activations or gradients reduce memory traffic when the recipe stores or communicates them in that format. Higher precision remains necessary where wide reductions, small parameter updates, or sensitive normalization can accumulate rounding error. The division is not universal. Frameworks choose dtypes per operation, kernels may accumulate in FP32, BF16 often avoids loss scaling, and optimizers may retain FP32 master weights and states. Mixed precision therefore coordinates formats across the training step instead of running every operation in one reduced precision.

8.6.3.1 Loss scaling

One of the key challenges with FP16 is its reduced dynamic range,²¹ which increases the likelihood of gradient values becoming too small to be represented accurately. Loss scaling addresses this issue by temporarily amplifying gradient values during backpropagation. Specifically, the loss value is scaled by a large factor (e.g., 2^{10}) before gradients are computed, ensuring they remain within the representable range of FP16.

Machine learning frameworks provide built-in support for mixed-precision training. PyTorch’s `torch.amp` (automatic mixed precision) library automates the process of selecting operation precision and can apply loss scaling when necessary.

8.6.3.2 Mixed-precision benefits

Mixed-precision benefits manifest across three dimensions that compound in practice:

- **Weight memory:** Decreases by 50 percent. A one-billion-parameter transformer’s weights require 4 GB in FP32 but 2 GB in FP16 or BF16. Total training-memory savings depend on retained higher-precision state.

²¹ | **FP16 (half-precision floating-point):** Its reduced normal range stems from using 5 exponent bits, compared with 8 in FP32; the smallest positive normal value is about 6.1×10^{-5} , but subnormal values extend to about 6×10^{-8} . Underflow behavior depends on the operation and hardware mode. Loss scaling reduces the chance that small gradients are rounded or flushed to zero.

- **Computational throughput:** Can increase when eligible operations use higher-throughput Tensor Core paths (section 8.6.3.3).
- **Communication bandwidth:** Can decrease when the distributed strategy communicates lower-precision tensors; some strategies still reduce or communicate gradients in FP32.

These benefits compound: a practitioner might simultaneously double batch size (memory savings), accelerate each iteration (Tensor Core throughput), and reduce gradient synchronization time (smaller tensors). On GPT-2, the combined effect is visible in both the memory budget and the training throughput.



Napkin Math 8.6: GPT-2 mixed precision memory savings

Problem: Can GPT-2 XL (1.5B parameters) be trained on a single V100 GPU, and what does mixed precision plus gradient checkpointing buy us in memory footprint?

FP32 baseline:

- Parameters and gradients (FP32): 12 GB (6 GB parameters + 6 GB gradients)
- Activations (batch = 4): ~71.7 GB
- Optimizer states (Adam m, v in FP32): 12 GB
- Total: ~95.7 GB (exceeds any single GPU)

FP16 mixed precision:

- Parameters (FP16): 1.5B parameters at 2 bytes each = 3 GB
- Activations (FP16): ~35.9 GB
- Gradients (FP16): 3 GB
- FP32 master weights: 6 GB (for precise optimizer updates)
- Optimizer states (Adam m, v in FP32): 12 GB
- Total: ~59.9 GB (still tight, but manageable with optimizations)

Mixed precision + illustrative fourfold activation reduction:

- Activations reduced to ~9 GB through selective recomputation
- Estimated tensor total: ~33 GB, compared with 32 GB of V100 capacity, before temporary workspaces and allocator overhead

Systems insight: Mixed precision halves the parameter and gradient tensors in this recipe; checkpointing then trades recomputation for lower activation storage. At batch size 4, the simplified tensor estimate approaches V100 capacity, so an actual run may still require additional memory reduction after accounting for workspaces and allocator behavior. At batch 32, the same estimate requires ~95.7 GB, motivating the scaling techniques later in this chapter.

Memory savings alone do not guarantee correct or stable training. FP16's limited dynamic range demands specific implementation choices to keep gradients representable and weight updates nonzero.

Three implementation details support conventional FP16 mixed precision:

- **Loss scaling:** Multiplies the loss before backpropagation so that small gradients are less likely to underflow. Dynamic scalers adjust the factor in response to detected nonfinite gradients rather than relying on one universal starting value or range.
- **FP32 master weights:** In recipes that use them, preserve updates that could round away if applied directly to FP16 weights.
- **Autocast policies:** Select lower or higher precision per operation. Reductions, normalization, and softmax commonly use higher-precision accumulation, but the exact policy depends on the framework and kernel.

With these stability mechanisms in place, the gains from mixed precision are measurable in throughput and dollars.

Mixed precision is not automatic. FP16's restricted exponent range can cause overflow or underflow, and subnormal handling varies by operation and hardware. Dynamic loss scaling addresses gradient-range failures, while monitoring for nonfinite losses, gradients, and activations remains essential. BF16 preserves FP32's exponent range and usually avoids loss scaling, but its shorter mantissa still changes rounding behavior. For small or non-Tensor-Core-dominated workloads, conversion and scaling overhead can outweigh the performance benefit; profiling, rather than a parameter-count threshold, determines the result.

Napkin Math 8.7: GPT-2 mixed precision throughput and cost

Problem: Given the following illustrative throughput and billing assumptions for a GPT-2 XL job, what speed and cost differences follow from FP16 mixed precision?

Assumed V100 throughput:

- FP32 throughput: ~90 samples/s
- FP16 throughput: ~220 samples/s
- Speedup: 220 samples/s ÷ 90 samples/s ≈ 2.4× faster training

Cost impact on a cluster of 32 GPUs:

- Assumed FP32 billing: \$50,000 for 2 weeks
- Assumed FP16 billing: \$28,000 for 1.2 weeks
- Wall-clock gain: 2 weeks ÷ 1.2 weeks ≈ 1.7×, below the 2.4× throughput speedup, since the assumed schedule includes work the FP16 kernels do not accelerate
- Time saved: 2 weeks – 1.2 weeks = 0.8 weeks (5.6 days)
- Cost saved: \$50,000 – \$28,000 = \$22,000

Quality gate: The cost comparison is valid only if a controlled validation run shows that the FP16 recipe meets the same acceptance criteria as the FP32 baseline. The scenario does not assume a particular perplexity difference.

Systems insight: Under these assumptions, the same cluster of 32 GPUs changes from 2 weeks and \$50,000 to 1.2 weeks and \$28,000. This is scenario arithmetic, not a benchmark claim; a real decision requires measured throughput, billing, and quality validation for the target workload.

8.6.3.3 Mixed-precision hardware support

Understanding how modern hardware implements reduced-precision arithmetic explains why mixed precision changes wall-clock time in addition to memory capacity. The performance gains from FP16 and BF16 computation come from specialized hardware units designed for low-precision tensor operations.²² Those units trade numerical range or precision for much higher matrix throughput, while the training recipe keeps numerically sensitive accumulations in safer formats.

NVIDIA introduced Tensor Cores in their Volta architecture (2017) as dedicated matrix multiplication units optimized for mixed-precision workloads. Unlike standard CUDA cores that process scalar or small vector operations, Tensor Cores perform 4×4 matrix multiply-accumulate operations in a single clock cycle. For FP16 inputs, a single Tensor Core executes:

$$\mathbf{D}_{\text{tc}} = \mathbf{A}_{\text{tc}} \times \mathbf{B}_{\text{tc}} + \mathbf{C}_{\text{acc}}$$

where $\mathbf{A}_{\text{tc}}$ and $\mathbf{B}_{\text{tc}}$ are FP16 inputs and $\mathbf{C}_{\text{acc}}$ and $\mathbf{D}_{\text{tc}}$ may use FP32 accumulation. Higher-precision accumulation reduces rounding error, although it does not eliminate numerical error or cancellation.

The peak numbers make the upper bound concrete. An NVIDIA A100 GPU exposes 19.5 TFLOP/s of FP32 throughput on standard CUDA cores and 312 TFLOP/s of FP16 or BF16 Tensor Core throughput, a 16× peak ratio. H100-class accelerators add FP8 Tensor Cores at 1,979 TFLOP/s without sparsity, about 2× the H100 FP16/BF16 Tensor Core rate. These figures are hardware ceilings, not end-to-end training promises: matrix multiplications map naturally to Tensor Cores, but data movement, non-Tensor-Core kernels, communication, loss scaling, and optimizer work remain. A transformer's attention mechanism computing $\mathbf{QK}^T$ for a tensor shaped by batch size B , number

²² | **Tensor Core:** A specialized matrix-multiply-accumulate unit that supports reduced-precision inputs and higher-precision accumulation. Supported tile shapes, alignment constraints, and peak-throughput ratios vary by architecture, dtype, sparsity mode, and library kernel.

of heads N_{heads} , sequence length S , and head dimension d_{head} requires $2 \times B \times N_{\text{heads}} \times S^2 \times d_{\text{head}}$ FLOPs; this portion can accelerate dramatically, while the rest of the step determines the realized speedup.

Brain Float 16 (BF16) maintains FP32's 8-bit exponent while reducing the mantissa to 7 bits. This design choice prioritizes dynamic range preservation over precision, which matters for gradient-based learning where values span many orders of magnitude. Google's TPUs natively support BF16, while NVIDIA's Ampere architecture (A100) and newer provide full hardware support.

The range advantage of BF16 over FP16 appears when gradients or intermediate values fall outside FP16's normal range. FP16's smallest positive normal value is 6.1×10^{-5} , while its subnormal floor is approximately 6×10^{-8} .²³ BF16's smallest positive normal value is approximately 10^{-38} , matching FP32's exponent range. BF16 therefore usually avoids the loss scaling needed by FP16, although accumulation precision and rounding still matter.

NVIDIA's Hopper architecture (H100) introduces FP8 support with two formats. E4M3 uses four exponent bits and three mantissa bits (prioritizing precision for forward pass weights and activations), while E5M2 uses five exponent bits and two mantissa bits (prioritizing dynamic range for backward pass gradients).

FP8 training doubles Tensor Core throughput again (1.98 PFLOP/s on H100 dense vs. 0.99 PFLOP/s for FP16 dense, without sparsity). However, FP8's severely limited precision requires per-tensor scaling factors maintained in higher precision, adding algorithmic complexity. Table 8.11 summarizes when each precision is appropriate:

Table 8.11: Precision Selection Guide: Format choice depends on hardware support, kernel coverage, numerical behavior, and validation results. FP8 requires an explicit scaling strategy; BF16 provides a wider exponent range than FP16; FP16 often uses loss scaling; and FP32 remains available for sensitive operations or unsupported lower-precision paths.

Precision	When to Use	Hardware Requirement
FP8	Throughput-oriented training with validated scaling	Supported FP8 hardware
BF16	Wide exponent range without FP16 loss scaling	Supported BF16 accelerator
FP16	Higher mantissa precision with managed loss scaling	Supported FP16 accelerator
FP32	Higher precision or unsupported lower-precision path	GPU

The decision begins with workload requirements and available hardware, then ends with empirical validation. No format is universally correct for an architecture class.

Reduced precision accelerates computation and alleviates memory bandwidth bottlenecks. Modern GPUs are increasingly compute-bound rather than bandwidth-bound for large matrix operations, but data movement still limits performance for smaller operations. A100's specifications illustrate this:

- HBM2e bandwidth: 2,039 GB/s
- FP32 throughput (19.5 TFLOP/s): worst-case demand is 78 TB/s if every FLOP needs new data
- Actual requirement (with data reuse): Much lower, but bandwidth-limited for operations with low arithmetic intensity

Using FP16 or BF16 halves the bytes per stored value relative to FP32, but realized memory traffic also depends on caching, fusion, accumulator formats, and extra conversions. Bandwidth-bound operations can therefore speed up even without Tensor Core arithmetic, although the gain is workload-dependent rather than an automatic doubling.

Modern frameworks abstract hardware complexity through automatic operation routing (chapter 7). The framework runtime determines which operations benefit from reduced precision and which require FP32 for numerical stability. PyTorch's automatic mixed precision manages precision selection and loss scaling transparently (listing 8.5).

The `autocast` context applies a framework-defined dtype policy to eligible operations; it does not imply that every matrix operation is low precision or that every normalization and loss operation is FP32. The example uses FP16 and a gradient scaler. A BF16 recipe generally omits gradient scaling.

²³ | **FP16 subnormal handling:** IEEE FP16 represents subnormals below 6.1×10^{-5} down to approximately 6×10^{-8} , but some accelerator operations or modes flush subnormal inputs or results to zero. The behavior must be checked for the target hardware and kernel.

Table 8.12 summarizes common options across GPU generations, but the final choice requires workload validation.

Listing 8.5: Mixed-Precision Training: Autocast selects eligible lower-precision operations, while gradient scaling is enabled for the FP16 path and omitted for BF16.

```
import torch

model = TransformerModel().cuda()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-4)
amp_dtype = (
    torch.float16
) # torch.bfloat16 usually does not need loss scaling
scaler = torch.amp.GradScaler(
    "cuda", enabled=(amp_dtype == torch.float16)
)

for inputs, targets in dataloader:
    optimizer.zero_grad()

    # Automatic precision selection per operation
    with torch.amp.autocast("cuda", dtype=amp_dtype):
        output = model(inputs)
        loss = criterion(output, targets)

    # Scale loss to prevent gradient underflow
    scaler.scale(loss).backward()

    # Unscale gradients before optimizer step
    scaler.step(optimizer)
    scaler.update() # Adjust scaling factor dynamically
```

Table 8.12: Precision Strategy by GPU Architecture: V100 supports FP16 Tensor Core training, A100 adds native BF16, and H100 adds FP8 paths whose use requires recipe-specific validation.

Architecture	Recommended Precision	Key Considerations
V100 (Volta)	FP16 with loss scaling	No native BF16 Tensor Core path; clipping depends on the training recipe
A100 (Ampere)	BF16 or FP16	BF16 avoids FP16 loss scaling; TF32 can accelerate eligible FP32 matrix operations
H100 (Hopper)	BF16, FP16, or validated FP8	FP8 requires an FP8-aware recipe; 1,979 TFLOP/s is a peak ceiling

An illustrative single-GPU GPT-2 scenario (1.5B) shows how assumed throughput might evolve with hardware and precision: 18 samples/s for V100 FP32, 45 samples/s for V100 FP16, 165 samples/s for A100 BF16, and 380 samples/s for H100 FP8. These values are scenario inputs rather than a controlled cross-generation benchmark, because software versions, kernels, batch shapes, and FP8 recipes also change. The valid systems lesson is that low-precision algorithms and specialized hardware must be evaluated together.

8.6.4 FlashAttention: IO-aware attention optimization

Mixed-precision training addresses two bottlenecks: compute throughput, because Tensor Cores operate faster on FP16, and memory capacity, because each value uses fewer bytes. For transformer models during training, however, a third bottleneck often dominates: memory bandwidth. The attention mechanism’s quadratic intermediate matrices must be repeatedly loaded and stored during the forward pass and accessed again during backpropagation. Even with reduced precision, the sheer volume of memory traffic can leave compute units idle while the processor waits for data.

FlashAttention²⁴ (T. Dao et al. 2022) addresses this bandwidth bottleneck by optimizing how data flows between memory hierarchies. Processing attention in SRAM-sized tiles avoids materializing

²⁴ | **FlashAttention:** Introduced by T. Dao et al. (2022), it preserves exact attention while reducing HBM traffic through SRAM-sized tiling; the key insight was to treat attention as an IO problem: avoid materializing the full $S \times S$ score matrix in HBM while retaining dense $\mathcal{O}(S^2)$ score computation. Performance depends on shape, hardware, and baseline. FlashAttention-2 (Dao 2023) further improved parallelism and work partitioning, reporting 50–73 percent of theoretical peak on A100.

the full $S \times S$ matrix in HBM and reduces traffic without changing the mathematical output. The original work reports up to $3\times$ speedups on evaluated workloads; gains depend on tensor shape, hardware, and the baseline implementation.

8.6.4.1 The standard attention memory bottleneck

Standard self-attention (chapter 6) computes relationships between all positions in a sequence. For an input sequence of length S , the mechanism computes an $S \times S$ attention matrix according to equation 8.16:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}} \right) \mathbf{V} \quad (8.16)$$

Here, $\mathbf{Q}$, $\mathbf{K}$, and $\mathbf{V}$ are the query, key, and value matrices for the sequence, and d_k is the key-vector dimension used to scale the dot products before softmax. The $\mathbf{Q}\mathbf{K}^T$ term is the source of the $S \times S$ score matrix, which is why attention becomes an I/O problem even when the arithmetic itself maps well to Tensor Cores.

The memory bottleneck emerges from materializing the $S \times S$ intermediate matrices for scores and probabilities. For a sequence length of 4,096 with embedding dimension 64 (typical for a single attention head), the attention score matrix alone requires $4,096^2 \times 4$ bytes = 67.1 MB in FP32. With 16 heads, this grows to 1.1 GB just for intermediate attention matrices, not including the keys, queries, values, or output tensors.

GPU memory hierarchy creates the opportunity for I/O-aware attention. HBM offers large off-chip capacity, while much smaller on-chip SRAM has substantially higher bandwidth and lower latency. Attention implementations that materialize the quadratic intermediates in HBM must write and later read them during backpropagation. The fraction of time spent on those transfers depends on sequence length, head dimension, dtype, hardware, and kernel implementation; it is not a fixed property of GPT-2-scale models.

The backward pass requires the information represented by the attention probabilities, but an implementation may either save those intermediates or recompute them. For $\mathbf{A}_{\text{attn}} = \text{softmax}(\mathbf{Z}_{\text{attn}})$ and $\mathbf{Z}_{\text{attn}} = \mathbf{Q}\mathbf{K}^T / \sqrt{d_k}$, one gradient expression is:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{Q}} = \frac{1}{\sqrt{d_k}} \cdot \text{dsoftmax} \left(\frac{\partial \mathcal{L}}{\partial \mathbf{A}_{\text{attn}}}, \mathbf{A}_{\text{attn}} \right) \cdot \mathbf{K}$$

where dsoftmax denotes the softmax Jacobian-vector product. A conventional implementation may save the probability matrix for this computation; saving both probabilities and raw scores is not mathematically required. FlashAttention instead saves compact normalization statistics and recomputes tiled score and probability blocks during backward propagation, avoiding full quadratic intermediates in HBM.

8.6.4.2 IO-aware attention through tiling

FlashAttention eliminates the need to materialize full $S \times S$ attention matrices in HBM by computing attention incrementally through tiling. Instead of computing the entire attention matrix at once, the algorithm partitions $\mathbf{Q}$, $\mathbf{K}$, and $\mathbf{V}$ into tiles small enough to fit in fast SRAM, computes attention scores for these tiles, and incrementally accumulates results.

The key algorithmic insight relies on the mathematical structure of softmax attention. Standard attention computes:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}} \right) \mathbf{V}$$

Algorithm 8.3 decomposes this computation by partitioning the queries and the keys/values into tiles small enough to fit in SRAM, then streaming over the key/value tiles while maintaining the running softmax statistics needed to assemble each output tile, without ever forming the full score matrix.

No full $S \times S$ score or probability matrix is written to HBM. A 128×128 FP32 score tile, for example, occupies 65.5 KB compared with 67.1 MB for a 4096×4096 matrix. Other inputs, outputs, and running statistics remain, so the tile is not necessarily the largest tensor in the entire kernel.

Full 67.1 MB

Tile 65.5 KB

log scale

FlashAttention swaps the full attention matrix for small SRAM tiles.

Algorithm 8.3 FlashAttention: tiled attention with online softmax**Require:** queries $\mathbf{Q}$, keys $\mathbf{K}$, values $\mathbf{V}$ for a length- S sequence; tile size b **Ensure:** attention output $\mathbf{Y}$, without materializing the $S \times S$ score matrix

```

1: for each query tile  $\mathbf{Q}_i$  do
2:   in SRAM:  $\mathbf{Y}_i^{\text{acc}} \leftarrow \mathbf{0}$ ,  $m_i \leftarrow -\infty$ ,  $l_i \leftarrow 0$ 
3:   for each key/value tile  $(\mathbf{K}_j, \mathbf{V}_j)$  do
4:     load  $\mathbf{Q}_i, \mathbf{K}_j, \mathbf{V}_j$  into SRAM;  $\mathbf{Z}_{ij} \leftarrow \mathbf{Q}_i \mathbf{K}_j^\top / \sqrt{d_k}$ 
5:      $m'_i \leftarrow \max(m_i, \text{rowmax}(\mathbf{Z}_{ij})) \triangleright \text{new stable max}$ 
6:      $\mathbf{Y}_i^{\text{acc}} \leftarrow \mathbf{Y}_i^{\text{acc}} \odot e^{m_i - m'_i} + e^{\mathbf{Z}_{ij} - m'_i} \mathbf{V}_j$ 
7:      $l_i \leftarrow l_i \odot e^{m_i - m'_i} + \text{rowsum}(e^{\mathbf{Z}_{ij} - m'_i})$ ;  $m_i \leftarrow m'_i$ 
8:     discard  $\mathbf{Z}_{ij} \triangleright \text{never written to HBM}$ 
9:   end for
10:   $\mathbf{Y}_i \leftarrow \mathbf{Y}_i^{\text{acc}} / l_i$ ; write  $\mathbf{Y}_i$  to HBM
11: end for

```

The online softmax algorithm enables this decomposition. Traditional softmax requires knowing all inputs before computing any output: $\text{softmax}(x)_i = e^{x_i} / \sum_j e^{x_j}$. FlashAttention uses an incremental formulation that updates softmax statistics as new blocks arrive, tracking the running maximum m (for numerical stability) and denominator l as each block is processed, then rescaling accumulated outputs accordingly.

8.6.4.3 Memory and IO complexity analysis

FlashAttention improves both activation memory and memory IO, which are the limiting costs in bandwidth-bound attention. Standard attention requires $\mathcal{O}(S^2)$ memory to store score matrices $\mathbf{Z}_{\text{attn}}$ and attention-probability matrices $\mathbf{A}_{\text{attn}}$ across all sequence positions. FlashAttention reduces the stored activation footprint to $\mathcal{O}(S)$ by keeping only input and output tensors $(\mathbf{Q}, \mathbf{K}, \mathbf{V}, \mathbf{Y})$ plus a small SRAM buffer for the current tile.

For $S = 4096$, $d_{\text{head}} = 64$: Standard attention requires $4096^2 \times 4 \text{ bytes} = 67.1 \text{ MB}$ per head. FlashAttention requires only $(3 \times 4096 \times 64) \times 4 \text{ bytes} \approx 3.1 \text{ MB}$ per head, a 21.3 $\times$ reduction.

The IO pattern changes for the same reason. Standard attention reads $\mathbf{Q}$, $\mathbf{K}$, and $\mathbf{V}$ from HBM, writes score, probability, and output matrices during the forward pass, then rereads those stored matrices during the backward pass before writing $d\mathbf{Q}$, $d\mathbf{K}$, and $d\mathbf{V}$. The resulting HBM traffic includes the same $\mathcal{O}(S^2)$ term that made the activation footprint large. FlashAttention forms score and probability blocks in SRAM and never writes the full $S \times S$ tensors to HBM. In the backward pass, it recomputes the needed blocks rather than reading stored full matrices. The exact HBM IO complexity depends on tile size and available SRAM because key and value blocks are streamed repeatedly across query blocks, so the bound is not simply $\mathcal{O}(S \cdot d)$.

For large sequence lengths, FlashAttention reduces HBM traffic by avoiding writes and rereads of the $S \times S$ score and probability matrices. With $S = 4096$ and $d_{\text{head}} = 64$, the activation memory footprint falls from hundreds of MB per head for stored attention matrices to a few MB for tiled inputs and outputs, while the precise bandwidth reduction depends on hardware and kernel tiling.

Both approaches require $\mathcal{O}(S^2 d)$ asymptotic FLOPs for dense attention. FlashAttention's backward pass recomputes tiled attention intermediates from saved inputs and normalization statistics rather than storing full score and probability matrices. The resulting reduction in HBM traffic can shift the kernel toward compute-bound execution and produce a net speedup, but the final bottleneck remains shape- and hardware-dependent.

8.6.4.4 Implementation and hardware utilization

FlashAttention's performance gains materialize through careful exploitation of GPU memory hierarchy. Modern frameworks can dispatch eligible attention operations among fused and fallback implementations based on hardware, inputs, and layouts; profiling verifies which path ran and whether it helped. Listing 8.6 contrasts standard and optimized attention implementations.

Listing 8.6: Attention Implementation Comparison: Standard attention materializes the full $S \times S$ matrix in HBM, while FlashAttention uses PyTorch’s optimized implementation or the dedicated flash-attn library.

```
import torch
import torch.nn.functional as F

# Standard attention (materializes n-by-n matrix)
def standard_attention(q, k, v):
    # q, k, v: [batch, heads, seq_len, head_dim]
    scores = torch.matmul(q, k.transpose(-2, -1)) / (
        q.size(-1) ** 0.5
    )
    attn = F.softmax(scores, dim=-1) # n-by-n matrix in HBM
    output = torch.matmul(attn, v)
    return output

# FlashAttention (no n-by-n materialization)
def flash_attention(q, k, v):
    # Can use FlashAttention when inputs and hardware are eligible
    output = F.scaled_dot_product_attention(q, k, v)
    return output

# Explicit FlashAttention 2 (flash-attn library)
from flash_attn import flash_attn_func

def flash_attn_2(q, k, v):
    # q, k, v: [batch, seq_len, heads, head_dim]
    # Different layout for optimized memory access
    output = flash_attn_func(q, k, v)
    return output
```

8.6.4.5 Benchmark results

The benefits of FlashAttention become concrete when measured on real hardware. T. Dao et al. (2022) reports end-to-end GPT-style training speedups and separate attention-kernel benchmarks showing that IO-aware attention reduces memory traffic and improves runtime. Table 8.13 uses an illustrative A100-style scenario to show the same systems pattern; its timings and memory values are representative chapter numbers, not values reported verbatim by T. Dao et al. (2022).

Table 8.13: FlashAttention Benchmark Comparison: Illustrative per-call timing and peak memory for standard attention vs. FlashAttention on a 40 GB A100-style configuration across sequence lengths. OOM marks configurations where standard attention exceeds the 40 GB memory budget; the 8192-token row also shows that standard attention would exceed 80 GB.

Sequence Length	Standard Forward	Flash Forward	Standard Backward	Flash Backward	Memory (Standard)	Memory (Flash)
512	12 ms	8 ms	35 ms	18 ms	4.2 GB	2.8 GB
2048	45 ms	15 ms	120 ms	35 ms	18 GB	6 GB
4096	OOM	32 ms	OOM	85 ms	> 40 GB	12 GB
8192	OOM	68 ms	OOM	180 ms	> 80 GB	24 GB

In this illustrative 40 GB A100-style scenario, standard attention runs out of memory beyond 2048 tokens, while FlashAttention fits sequences up to 8192 tokens. Even at 2048 tokens where both fit, FlashAttention achieves 3× forward pass speedup and 3.4× backward pass speedup.

Subsequent versions have continued improving performance: FlashAttention-2 (Dao 2023) achieved 1.5–2× additional speedup through better parallelism and register allocation, while FlashAttention-3 (Shah et al. 2024) exploits FP8 tensor cores and asynchronous memory operations

on Hopper GPUs and reports reaching approximately 740 TFLOP/s on H100, around 75 percent of theoretical peak.

8.6.4.6 When to use FlashAttention

For transformer training with long sequences, FlashAttention is usually the first attention implementation to test. It often becomes essential once sequence length pushes standard attention into an HBM-capacity or HBM-bandwidth bottleneck, especially around the multi-thousand-token regime where the $S \times S$ attention matrices dominate memory. A100- and H100-class GPUs with fast on-chip SRAM benefit most. The returns diminish for short sequences where tiling overhead is not worthwhile, on older GPU architectures without comparable SRAM bandwidth, and for nonattention architectures such as convolutional neural networks and multilayer perceptrons.

In practice, deep learning frameworks handle much of FlashAttention’s integration, but the decision still depends on layout, precision, and memory budget. PyTorch’s `scaled_dot_product_attention` can dispatch to FlashAttention when hardware, inputs, and backend constraints permit; hand-written attention must use an eligible primitive. Three practical checks determine whether FlashAttention is appropriate:

1. Ensure tensor layouts match library expectations (contiguous memory, correct dimension ordering)
2. Use FP16 or BF16 for maximum speedup (FlashAttention optimized for mixed precision)
3. Combine with gradient checkpointing when additional activation-memory savings are needed.

The integration is often a single-line change—swapping a manual attention call for `F.scaled_dot_product_attention` (see listing 8.6). The engineering decision is still the same one established by the roofline analysis: use the optimized primitive when the bottleneck is HBM traffic, and let the library manage the tiling and SRAM scheduling needed to remove it.

8.6.4.7 Systems implications and broader principles

FlashAttention exemplifies a fundamental systems engineering principle: **IO-aware algorithm design**. Many accelerators are compute-abundant relative to memory bandwidth. For a bandwidth-bound kernel, traffic can govern runtime even when FLOPs do not change.

This principle extends beyond attention. In IO-aware matrix multiplication, tiling algorithms like those in CUTLASS minimize DRAM traffic by maximizing data reuse in fast caches. A naive $m \times m$ matrix multiply performs $\mathcal{O}(m^3)$ FLOPs with $\mathcal{O}(m^2)$ memory traffic, while blocked algorithms maintain $\mathcal{O}(m^3)$ FLOPs but reduce cache misses through locality optimization.

The same byte-movement principle reappears in communication-efficient distributed training, where gradient compression trades extra computation (compression/decompression) for reduced network bandwidth consumption. Low-power edge devices with limited memory bandwidth benefit even more from IO-aware algorithms. Trading 10 percent more arithmetic to halve memory traffic approaches a $2\times$ energy reduction in the limit where data movement accounts for the entire budget, and falls short of that bound whenever arithmetic carries a meaningful share, because halving one additive term of a sum cannot more than halve the sum. This section establishes the IO-aware heuristic; distributed and edge deployment settings reuse the same byte-movement logic.

FlashAttention transforms practical model training capabilities. By avoiding materialization of the $\mathcal{O}(S^2)$ attention matrices, it can enable three capabilities:

- **Longer sequences:** In this illustrative configuration, enables $4\times$ context length on the same hardware, moving GPT-2 on A100 from 2K to 8K context.
- **Larger batch sizes:** In the same configuration, doubles batch size through freed memory, potentially improving utilization; convergence effects still require validation.
- **Deeper models:** Reduces activation memory so more layers fit in the same memory budget.

In the chapter’s attention-only scenario, the optimized kernel shifts the context boundary from 2K to 8K tokens; full 7-billion-parameter training still depends on state, sharding, batch size, and other activations.

The technique demonstrates that systems-level algorithms can exploit memory hierarchy for large activation-memory reductions and workload-dependent speedups beyond hardware scaling alone.

Treating memory bandwidth as the primary constraint and compute as abundant is a recurring pattern in ML performance optimization.

FlashAttention addresses memory bandwidth bottlenecks during computation, but another class of memory constraints exists: the sheer capacity required to store activations and optimizer states simultaneously. When models or batch sizes exceed GPU memory capacity, two complementary techniques trade computation for memory.

8.6.5 Gradient accumulation and checkpointing

Training large models requires substantial memory for activations, gradients, and parameters. When memory constrains batch size or model complexity, gradient accumulation²⁵ forms a larger effective batch from sequential micro-batches, while checkpointing discards and recomputes selected activations. These techniques become standard tools when memory is the binding constraint.

8.6.5.1 Gradient accumulation and checkpointing mechanics

Gradient accumulation and activation checkpointing operate on distinct principles, but both aim to optimize memory usage during training by modifying how forward and backward computations are handled. Gradient accumulation changes how many examples contribute to one update; checkpointing changes which activations remain resident between the forward and backward passes.

Gradient accumulation To execute large batch training on memory-constrained accelerators, systems split effective batches across smaller sequential micro-batches. In figure 8.13, trace how micro-batch gradient vectors $\delta_1, \delta_2, \delta_3$ accumulate into a single unified gradient tensor before triggering an optimizer update.




Figure 8.13: Gradient Accumulation: Three micro-batches each compute independent losses and gradients, which sum into a single combined gradient for one parameter update. This simulates training with a batch three times larger without requiring the memory to hold all samples simultaneously.

In PyTorch, gradient accumulation is usually implemented by dividing each micro-batch loss by the number of accumulation steps and calling `optimizer.step()` only after processing the entire effective batch. Under that averaging convention, no learning-rate adjustment is needed solely because of accumulation. The implementation follows five steps:

1. Perform the forward pass for a micro-batch.
2. Compute the gradients during the backward pass.
3. Accumulate the gradients into a buffer without updating the model parameters.
4. Repeat steps 1–3 for all micro-batches in the effective batch.
5. Update the model parameters using the accumulated gradients after all micro-batches are processed.

Gradient accumulation can reproduce the gradient of a larger batch under controlled assumptions. For an effective batch size $B = k \times b$ where k is the number of accumulation steps and b is the micro-batch size, equation 8.17 confirms that the accumulated gradient equals the true batch gradient:

$$\nabla \mathcal{L}_B = \frac{1}{B} \sum_{i=1}^B \nabla \mathcal{L}_i = \frac{1}{k} \sum_{j=1}^k \left(\frac{1}{b} \sum_{i \in \mathcal{B}_j} \nabla \mathcal{L}_i \right) \quad (8.17)$$

²⁵ | **Gradient accumulation:** Sums or averages gradients across k micro-batches before one optimizer step. It reduces resident activation memory relative to processing the same effective batch at once, while preserving total example-level arithmetic. Runtime and numerical equivalence depend on micro-batch efficiency, synchronization, stateful layers, stochastic operations, precision, and loss normalization.

This equivalence holds when each example's computation is identical and the loss is averaged consistently. The right-hand side shows that averaging k micro-batch gradients (each computed over b examples) produces the same result as computing the gradient over all $B = kb$ examples at once. Stateful layers such as BatchNorm, stochastic operations such as dropout and data augmentation, finite-precision accumulation order, and step-indexed schedules can break exact equivalence, so large-batch accumulation still requires validation.

Gradient accumulation exchanges memory capacity for computation time. Table 8.14 separates the memory benefit from the compute and scheduling costs that remain.

Table 8.14: Gradient Accumulation Trade-Offs: Accumulation reduces resident activation memory and preserves total example-level arithmetic. Wall-clock cost depends on micro-batch efficiency, kernel-launch overhead, and whether distributed synchronization is deferred until the accumulated update.

Dimension	Effect
Memory	$\mathcal{O}(b)$ instead of $\mathcal{O}(B)$, yielding a $k\times$ reduction in activation memory
Computation	Unchanged total FLOPs, as all B examples are still processed
Time	k micro-batch forward and backward passes precede each optimizer step; smaller micro-batches can reduce utilization and add launch overhead

Accumulation preserves the total example-level arithmetic but can change efficiency because smaller micro-batches use the accelerator differently and launch kernels more often. With distributed `no_sync()` accumulation, one gradient synchronization occurs after the k local micro-batches rather than after each one. The model in equation 8.18 gives the per-update time:

$$T_{\text{effective}} = \sum_{j=1}^k T_{\text{micro},j} + T_{\text{sync}} + T_{\text{update}} \quad (8.18)$$

Here $T_{\text{micro},j}$ is the measured forward/backward time for micro-batch j , T_{sync} is the synchronization time paid for the accumulated update, and T_{update} is the optimizer-step time.

The wall-clock penalty relative to processing the effective batch at once is workload-dependent. It must be measured rather than inferred from the memory-reduction factor.

When gradient accumulation is combined with distributed data parallelism across multiple machines, additional considerations arise for gradient synchronization timing and effective batch size calculation across the cluster. Advanced distributed systems texts treat these patterns in depth.

Activation checkpointing Activation checkpointing reduces memory usage during the backward pass by discarding and selectively recomputing activations. In standard training, activations from the forward pass are stored in memory for use in gradient computations during backpropagation. However, these activations can consume gigabytes of memory, particularly in deep networks.

Activation memory can be reduced by trading FLOPs for memory bandwidth during backpropagation. Contrast the forward pass in the top row of figure 8.14 with the backward pass in the bottom row, observing how intermediate activations are discarded and recomputed on demand from solid green checkpoint tensors.




Figure 8.14: Activation Checkpointing: Trading memory usage for recomputation during backpropagation enables training deeper neural networks. By storing only a subset of activations from the forward pass and recomputing others on demand, this technique reduces peak memory requirements at the cost of increased training time.

The implementation keeps the capacity trade-off explicit. The system splits the model into segments, retains activations only at segment boundaries during the forward pass, and recomputes intermediate activations during the backward pass when needed.

Frameworks like PyTorch provide tools such as `torch.utils.checkpoint` to simplify this process, but the recompute cost remains. Checkpointing is most effective for deep architectures with dozens or hundreds of layers, such as transformers or large convolutional networks, where activation storage exceeds the GPU’s capacity before arithmetic throughput is exhausted.

The synergy between gradient accumulation and checkpointing enables training of larger, more complex models. Gradient accumulation manages memory constraints related to batch size, while checkpointing optimizes memory usage for intermediate activations. Together, these techniques expand the range of models that can be trained on available hardware.

8.6.5.2 Optimal checkpoint placement strategy

For a network with N_L layers, each storing A bytes of activations, table 8.15 quantifies how the number and placement of checkpoints determines the memory-compute trade-off. Sub-linear checkpointing strategies can reduce memory consumption from $\mathcal{O}(N_L)$ to $\mathcal{O}(\sqrt{N_L})$ with only a fractional increase in total compute time, enabling the training of much deeper models on existing hardware.

Table 8.15: Checkpointing Memory-Compute Trade-Offs: Different checkpoint strategies trade memory savings against recomputation overhead. The optimal number of checkpoints balances these factors.

Strategy	Memory Cost	Recompute Cost
No checkpointing	$N_L \times A$	0 forward ops
Retain one boundary	A	$(N_L - 1)$ forward ops
k checkpoints	$k \times A + (N_L/k) \times A$	$(N_L - k)$ forward ops

If we set the derivative of total memory cost ($k \times A + (N_L/k) \times A$) to zero, we locate the optimum at $k_{\text{optimal}} = \sqrt{N_L}$. In the common $\sqrt{N_L}$ checkpointing scheme, this gives roughly one-third extra forward compute in exchange for the minimum memory footprint under this simplified model. For GPT-2 with forty-eight transformer layers, the contrast is stark: without checkpointing, memory equals $48 \times A$ (full activation storage). Optimal checkpointing ($\sqrt{48}$ approximately equals seven checkpoints) requires memory of $7 \times A + (48/7) \times A$ approximately equals $14 \times A$, achieving 71 percent memory savings with approximately 33 percent compute overhead.

Not all operations are equally expensive to recompute, which motivates selective checkpointing. Table 8.16 compares where activation memory is large enough to justify recomputation.

Table 8.16: Selective Checkpoint Placement: Placement compares the state saved at a block boundary with the work required to reconstruct it. Attention and feed-forward costs depend on sequence length, width, kernel implementation, and whether memory-efficient attention is already in use.

Layer type	Memory cost	Recompute cost	Checkpointing strategy
Attention blocks	Shape-dependent; materialized attention can add quadratic intermediates	High and sequence-length-dependent	Checkpoint block boundaries when saved state justifies recomputation
Feed-forward blocks	Often dominated by expanded hidden activations	High and width-dependent	Checkpoint when expanded activations dominate memory
Normalization	Small standalone outputs	Low	Usually handled within a larger checkpointed block

The practical rule is to spend recomputation on high-memory layers, not on operations whose activations are already cheap to retain.

8.6.5.3 Memory and computational benefits

Gradient accumulation produces an effective batch larger than the resident micro-batch without storing all of its activations simultaneously. A larger effective batch reduces sampling variance but does not guarantee faster convergence or better generalization. The technique is particularly valuable when the desired effective batch cannot fit in memory at once.

Activation checkpointing significantly reduces the memory footprint of intermediate activations during the forward pass, allowing training of deeper models. By discarding and recomputing activations as needed, checkpointing frees up memory for larger models, additional layers, or higher resolution data. This trade-off is critical in architectures like transformers that require substantial memory for intermediate computations.

Both techniques improve scalability by reducing the memory required before adding hardware. Returning to lighthouse 8.1, gradient accumulation is essential for achieving the target batch size within V100 memory constraints.

Napkin Math 8.8: GPT-2 gradient accumulation strategy

This illustrative GPT-2-scale scenario uses assumed prices and durations to show gradient accumulation; it does not reconstruct GPT-2's historical training bill.

Memory Constraints

- With the chapter's simplified memory recipe, a V100 with 32 GB holds a per-GPU micro-batch of $B = 4$.
- A resident effective batch of $B = 512$ would therefore require $512 \div 4 = 128$ GPUs.

Accumulation Configuration

- Use 8 GPUs and accumulate 16 micro-batches of 4 examples each.
- Each accelerator processes $4 \times 16 = 64$ examples per update, giving a global effective batch of $8 \text{ GPUs} \times 64 = 512$.

A distributed `no_sync()` context can defer AllReduce until the final micro-batch.

Trade-off Analysis

- Idealized wall-clock time: $16\times$ longer because fewer accelerators perform the same arithmetic work
- Memory overhead: accumulation normally reuses parameter-gradient buffers, although bookkeeping and temporary storage remain
- Communication: both configurations synchronize once per 512-sample update, but participant count and topology differ
- Equal-work cost: \$688K either way, before utilization and communication effects

Why this works: Under the equivalence conditions stated above, averaged per-example gradients are additive. Each accelerator first accumulates a local gradient over its 64 examples:

$$\frac{1}{16} \sum_{j=1}^{16} \left[\frac{1}{4} \sum_{k=1}^4 \nabla \mathcal{L}(x_{jk}) \right]$$

AllReduce then averages those local gradients across 8 GPUs, producing the global effective batch:

$$\frac{1}{8} \sum_{g=1}^8 \nabla \mathcal{L}_{\text{local}}^{(g)}, \quad B_{\text{global}} = 8 \times 64 = 512$$

Operational comparison

Resident configuration:

- $128 \text{ GPUs} \times 4 \text{ examples} = 512$ resident examples per update
- Gradient synchronization spans 128 GPUs
- Hourly cost: $128 \text{ GPUs} \times \$16/\text{hour} = \$2,048/\text{hour}$

Accumulated configuration:

- $8 \text{ GPUs} \times (4 \text{ examples} \times 16 \text{ micro-batches}) = 512$ examples per update

- Gradient synchronization occurs once per update across 8 GPUs
- Hourly cost: 8 GPUs $\times$ \$16/hour = \$128/hour
- Hourly reduction: \$1,920/hour, or 93.8 percent

Quality gate

- Match the examples and their order, loss normalization, optimizer state and step schedule, and stochastic-layer behavior.
- Compare held-out loss or perplexity and the optimization trajectory across repeated runs rather than assuming that the effective batch alone guarantees equivalent quality.
- Expect small round-off differences from the changed summation order; require convergence within a declared tolerance rather than bitwise identity. No perplexity result is assumed here.

Systems insight: Accumulation lowers accelerator count and hourly burn, but ideal scaling preserves accelerator-hours; actual cost depends on utilization and communication.

Listing 8.7 places one AllReduce after sixteen local micro-batches, making the synchronization boundary explicit.

Listing 8.7: Gradient Accumulation Training Loop: Sixteen local micro-batches accumulate before one AllReduce and optimizer update.

```
optimizer.zero_grad()
for step in range(16): # Accumulation steps
    micro_batch = next(dataloader) # 4 samples
    loss = model(micro_batch) / 16 # Average effective-batch loss
    loss.backward() # Accumulate gradients
# Now gradients represent 64 local samples
all_reduce(gradients, op=SUM) # Sum across 8 GPUs
gradients /= 8 # Convert the sum to the global mean
optimizer.step() # Update with effective batch=512
```

8.6.5.4 Practical considerations

Gradient accumulation is most valuable when optimal batch sizes exceed GPU memory capacity. Transformer language models often use effective batches of hundreds of thousands to millions of tokens; if memory only permits a smaller number of sequences per device, accumulation bridges the gap without requiring additional hardware. Activation checkpointing complements this by enabling deeper architectures and is commonly used in transformer-scale training, as well as in activation-heavy dual-network configurations such as generative adversarial networks.

Both techniques introduce explicit trade-offs. Activation checkpointing recomputes discarded activations during backward propagation; with ordinary segment checkpointing, each discarded activation is typically reconstructed once, while recursive schedules may behave differently. Gradient accumulation reduces parameter-update frequency because each update follows k micro-batches. When each micro-batch loss is divided by k , the accumulated gradient follows the effective-batch mean convention. If losses are summed instead, the gradients must be rescaled before the optimizer step or the learning-rate convention changed accordingly. Framework and codebase conventions differ, making normalization a common source of subtle bugs.

8.6.6 Optimization technique comparison

Table 8.17 synthesizes three of the four core optimization strategies, contrasting their primary goals, mechanisms, and trade-offs. FlashAttention complements them by addressing attention’s memory-bandwidth bottleneck through IO-aware tiling, reducing auxiliary attention memory from $\mathcal{O}(S^2)$ to $\mathcal{O}(S)$ and reporting speedups of up to $3\times$ on evaluated workloads. Selecting an appropriate strategy depends on the bottleneck identified through profiling.

These techniques target different constraints. Prefetching addresses data starvation, mixed precision can accelerate eligible computation and reduce tensor storage, FlashAttention reduces attention memory traffic, and gradient accumulation enables effective batches that would otherwise exceed activation memory. Their benefit depends on the measured bottleneck, and applying one does not guarantee that it removes the end-to-end constraint.

Table 8.17: Optimization Strategies: The table compares prefetching, mixed precision, and the combined use of gradient accumulation and checkpointing. The first targets input stalls, the second changes storage and eligible arithmetic, and the third trades smaller resident activation sets for micro-batch or recomputation overhead.

Aspect	Prefetching and Overlapping	Mixed-Precision Training	Gradient Accumulation and Checkpointing
Primary Goal	Minimize data transfer delays and maximize accelerator utilization	Reduce tensor storage and accelerate eligible lower-precision operations	Overcome memory limitations during backpropagation and parameter updates
Key Mechanism	Asynchronous data loading and parallel processing	Combining lower- and higher-precision computation	Simulating larger batch sizes and selective activation storage
Memory Impact	Increases memory usage for prefetch buffer	Can reduce weights, activations, and communicated tensors	Reduces resident activations; parameter-gradient size is unchanged
Computation Speed	Improves by reducing idle time	Can accelerate eligible FP16/BF16 operations on supported hardware	May slow down due to recomputations in checkpointing
Scalability	Highly scalable, especially for large datasets	Enables training of larger models	Allows training deeper models on limited hardware
Hardware Requirements	Benefits from fast storage and multi-core CPUs	Supported lower-precision accelerator paths	Works on standard hardware
Implementation Complexity	Moderate (requires tuning of prefetch parameters)	Low to moderate (with framework support)	Moderate (requires careful segmentation and accumulation)
Main Benefits	Reduces training time, improves hardware utilization	Faster training, larger models, reduced memory usage	Enables larger batch sizes and deeper models
Primary Challenges	Tuning buffer sizes, increased memory usage	Potential numerical instability; FP16 may require loss scaling	Increased computational overhead, slower parameter updates
Ideal Use Cases	Large datasets, complex preprocessing	Large-scale models, especially in NLP and computer vision	Memory-constrained batches or activation-heavy models

The table compares techniques in isolation; the GPT-2 walkthrough shows how those techniques combine when one model must fit in memory, run fast enough, and stay within a realistic cost envelope.

8.6.7 GPT-2 optimization walkthrough

The memory-feasibility path for GPT-2 (1.5 billion parameters) training on a single V100 GPU establishes how these optimizations combine before extending the time, energy, and cost analysis to a 32-V100 training run.

 Napkin Math 8.9: GPT-2 optimization on V100

- Step 1: Establish the baseline memory footprint.
- The trace starts with GPT-2 XL with 1.5 billion parameters, batch size 4, sequence length 1024, FP32 tensors throughout, and a single-threaded synchronous data loader. This is a deliberately constrained configuration: small enough to discuss on one V100, but large enough that the first failure is immediate and unambiguous. This walkthrough uses batch size 4, the same configuration as the mixed-precision callout (napkin math 8.6); table 8.19 reports the batch-32 equivalents (597.8 GB FP32 baseline, 95.7 GB optimized), which exceed a single GPU even after optimization. At batch size 4, the FP32 baseline needs 95.7 GB, so the run fails the machine constraint on a 32 GB V100 before throughput matters.
- Step 2: Apply mixed precision.

The first intervention targets memory feasibility, not speed. Mixed precision changes the first term in the memory budget and reduces the footprint to 59.9 GB, a 37.5 percent reduction, but it still does not fit.

Step 3: Add gradient checkpointing.

Gradient checkpointing changes a different term by storing fewer activations and recomputing them during backpropagation. The scenario assumes a 4× activation-memory reduction, bringing the simplified tensor total to 33 GB. The 33 percent compute overhead follows the simplified equal-layer model; an actual run must measure both memory and time, including temporary workspaces.

Step 4: Examine an illustrative throughput profile.

Fitting the model is only the first diagnostic pass. Suppose a subsequent profile reports the following scenario values:

- Accelerator utilization: 45 percent
- Data loading: 40 percent of iteration time
- Compute: 35 percent of iteration time
- Memory transfers: 25 percent of iteration time

These assumed measurements would move the next investigation from machine capacity to the data axis in D-A-M terms, so additional memory optimization would not be the first wall-clock intervention.

Step 5: Apply prefetching and data pipeline optimization.

The candidate fix is to overlap input preparation with accelerator execution. In this scenario, configuring the data loader with eight workers, pinned memory, and a prefetch factor of 2 is assumed to raise accelerator utilization to 85 percent. A real profile must confirm the improvement and identify the remaining overhead.

Table 8.18 contrasts the naive baseline against the optimized configuration:

Table 8.18: Illustrative GPT-2 Optimization Profile: The batch-4 memory figures are calculated by the walkthrough; utilization, throughput, and epoch time are assumed scenario values rather than benchmark results. Table 8.19 reports batch-32 scenario totals.

Metric	Naive	Optimized	Improvement
Memory	95.7 GB	33 GB	2.9× reduction
Accelerator utilization	N/A	85%	Trainable
Throughput	N/A	1,200 tokens/s	—
Time per epoch	N/A	8.3 hours	—

Systems insight: If re-profiling confirms that capacity and input stalls are no longer binding, the next optimization decision should follow the newly observed bottleneck rather than assume that the run is compute-bound.

The case study illustrates why training optimization is an iterative diagnostic discipline rather than a menu of tricks. Each intervention answers the bottleneck exposed by the previous measurement: mixed precision reduces tensor storage but does not solve capacity alone, checkpointing accepts 33 percent additional compute to gain 2.9× memory reduction, and prefetching matters only after the model can run and profiling reveals data starvation. The systematic loop of profile, identify the bottleneck, apply a targeted technique, and re-profile transforms optimization from trial-and-error into engineering practice.

8.6.8 Optimization impact summary

The GPT-2 case study demonstrates how targeted techniques can compose. Its memory values follow the chapter's tensor model, while its time, power, electricity-price, and grid-carbon inputs define an illustrative operating scenario rather than a measured GPT-2 run.

Table 8.19 compiles the results implied by those cluster-level assumptions.

Table 8.19: Illustrative GPT-2 Training Scenario: The tensor model gives the stated memory totals. Training duration, constant cluster power, PUE, electricity price, and grid carbon intensity are scenario assumptions; under those assumptions, shorter runtime reduces energy and location-based emissions proportionally.

Metric	FP32 Baseline	Optimized	Technique Applied
Parameters	6 GB	3 GB	Mixed precision (FP16)
Gradients	6 GB	3 GB	Mixed precision (FP16)
Master Weights	0 GB	6 GB	AMP Overhead
Optimizer State (Adam)	12 GB	12 GB	Unchanged (FP32 moments)
Activations (batch=32)	573.8 GB	71.7 GB	Gradient checkpointing + FP16
Total Memory	597.8 GB	95.7 GB	—
Training Time (32 V100s)	14 days	8.4 days	Data prefetching + utilization
Energy Consumption	3,914 kWh	2,348 kWh	Assumed constant cluster power over each duration
Electricity Cost (\$0.10/kWh)	\$391.4	\$234.8	—
Carbon Footprint	1.7 t	1.0 t	Regional grid average (0.43 kg/kWh)

As table 8.19 shows, the assumed inputs produce a 6× memory ratio, a 1.7× wall-clock ratio, and a 40 percent energy reduction under constant modeled power. In practice, energy and cost must be computed from measured runtime and power because higher utilization can change the cluster’s power draw.

The single-machine optimization toolkit needed for this chapter is now exhausted. Mixed precision raises attainable Tensor Core throughput. FlashAttention reduces HBM traffic. Gradient checkpointing trades compute for memory. Prefetching can hide data-loading latency. Together these methods can train GPT-2-scale models much faster than an untuned baseline, but the chapter’s stated compute budget still cannot finish on one contemporary GPU in days.

Some models will not fit on one device, and some training runs remain impractically long after single-machine optimization. When training still exceeds acceptable time or memory budgets, the next option is to spread the computation across multiple devices. This transition introduces new bottlenecks, including communication overhead, synchronization costs, and fault tolerance requirements.


8.7 Scaling Training Systems

The GPT-2 walkthrough models how mixed precision and checkpointing reduce a batch-4 tensor budget to 33 GB, near a V100’s 32 GB capacity before workspaces and allocator overhead. At batch 32, the modeled total falls from 597.8 GB to 95.7 GB but still exceeds that V100. Larger-memory accelerators may move the boundary. A 70-billion-parameter model, however, requires about 140 GB for FP16 weights alone and substantially more for full-parameter training state, motivating sharding, offload, or multiple devices.

When a chosen workload still exceeds one device’s memory or schedule, spreading it across devices provides aggregate memory and compute. Other responses include reducing the model, batch, sequence length, or training scope; offloading state; and using parameter-efficient adaptation. Scaling beyond one device begins with multi-GPU configurations inside one machine and may extend across nodes. The major parallelism strategies make different communication trade-offs; a full treatment of distributed-training implementation lies beyond this chapter.

Not all workloads benefit equally from adding GPUs. The relationship among useful computation, communication, imbalance, and synchronization determines scaling efficiency. Figure 8.15 models this trade-off using an Amdahl-style formulation where non-scaling overhead creates a binding throughput ceiling. The widening gap between ideal linear speedup and achieved speedup visualizes the communication tax paid at scale.

The red curve’s widening gap from the dashed ideal is the communication tax made visible; shrinking that gap is what the distributed strategies in section 8.7.1 exist to do.






Figure 8.15: Communication Tax: An Amdahl-style model compares ideal linear scaling against workloads with varying non-scaling communication and synchronization fractions. Real communication overhead scales dynamically with device count, network topology, message size, and overlap efficiency.

8.7.1 Single-node multi-GPU training

Multi-GPU training within a single node, the scope of this book, predates large-scale distributed systems. AlexNet²⁶ (Krizhevsky et al. 2012) famously split its model across two GTX 580 GPUs because the 3 GB memory limit made the network too large to train on one GPU. The split also made cross-GPU connectivity an architectural choice. This single-node, multi-GPU configuration remains a useful entry point because it introduces the core parallelism strategies without the complexity of network communication.

The two foundational strategies, data parallelism and model parallelism, represent fundamentally different partitioning choices: data parallelism replicates the model and partitions data, while model parallelism partitions model state or computation. This distinction determines memory requirements, communication patterns, and scaling behavior.

8.7.1.1 Data parallelism

Data parallelism replicates the entire model on each GPU, with each processing different batches. After computing gradients locally, GPUs synchronize via gradient averaging. Figure 8.16 shows the data flow: input data splits into nonoverlapping batches, each accelerator computes forward and backward passes independently, then gradients aggregate before updating the shared model.

Data parallelism's appeal lies in its simplicity and efficiency. Each GPU runs the same forward-backward computation on different data, and replicas agree on gradients before updating parameters. Modern implementations bucket and overlap reductions with backpropagation. This makes data parallelism the natural first strategy when models fit in GPU memory. PyTorch's `DistributedDataParallel` and TensorFlow's `MirroredStrategy` automate synchronization, keeping the programming model close to single-GPU training despite per-iteration communication.

Unsharded data parallelism has a hard constraint: every GPU holds a complete model replica and ordinarily its local optimizer state. A 7-billion-parameter model uses 14 GB for FP16 weights alone, before gradients, optimizer states, or activations. Sharded data parallelism (such as the Zero Redundancy Optimizer, or ZeRO) distributes that state across devices: ZeRO-1 shards optimizer states, ZeRO-2 shards gradients, and ZeRO-3 (or FSDP) shards model parameters; otherwise, a model that exceeds local memory requires pipeline or tensor partitioning.

8.7.1.2 Model parallelism

Model parallelism partitions the model itself across GPUs, which becomes necessary when the model exceeds single-GPU memory. AlexNet used a simple form: certain layers resided on GPU 1, others on GPU 2, with activations passing between them. Figure 8.17 shows the forward and

²⁶ | **AlexNet (2012):** The paper reports 60M parameters and a 3 GB single-GPU ceiling, but does not attribute the limit to one memory component. Its model-parallel design placed roughly half the kernels on each GPU and allowed communication only in selected layers, explicitly trading memory capacity against communication.




Figure 8.16: Data Parallelism: Each GPU holds a complete model copy, processes a different data shard, and synchronizes gradients. Throughput can approach linear scaling only while computation dominates synchronization, input, and imbalance costs.

backward paths: data moves through model partitions on different devices, with gradients flowing backward during training.




Figure 8.17: Model Parallelism: The model is partitioned across devices, with intermediate activations passing between them. This enables training models larger than single-GPU memory at the cost of sequential dependencies.

 Systems Perspective 8.7: Data vs. model parallelism

When a selected workload is partitioned across N_{GPU} devices, the split may divide examples, model state, layers, or tensor operations. For a model with parameters P and batch size B , data and model parallelism provide two starting points.

Data parallelism (split the batch): For a fixed global batch, each of N_{GPU} workers processes about $1/N_{\text{GPU}}$ of the examples, but every worker holds a model replica and ordinarily its associated training state. Dense gradient communication is proportional to P per update. Scaling departs from the ideal as local batches shrink or synchronization and input costs grow.

Model parallelism (split model computation): Each accelerator stores and computes a portion of the model. An even partition may approach P/N_{GPU} parameters per worker, but embeddings, buffers, imbalance, and replicated state alter that estimate. Communication depends on the partition: layerwise pipeline parallelism exchanges activations at stage boundaries, while tensor parallelism performs collectives within partitioned layers.

Systems insight: Data parallelism is the simplest candidate when the complete training state fits per worker and more throughput is needed. Model-state, pipeline, or tensor partitioning becomes necessary when it does not, often in combination with data parallelism.

When a model exceeds the memory capacity of a single accelerator, its depth can be partitioned across sequential devices. Examine the pipeline stage assignment in figure 8.18, observing how contiguous groups of transformer blocks map onto distinct devices across the execution chain.




Figure 8.18: Layer-Wise Model Partitioning: Distributing consecutive transformer layers across sequential devices bounds network communication to intermediate activations (forward) and boundary gradients (backward). However, sequential layer dependencies leave downstream devices idle during forward passes unless micro-batches are pipelined.

Model parallelism’s challenge is idle time. While Device 1 computes layers 1–6, Devices 2–4 sit idle waiting for activations. During the backward pass, the problem reverses: Device 4 computes first while others wait. This “pipeline bubble” means naive model parallelism can waste much of the available accelerator time even with careful partitioning. Pipeline parallelism addresses this inefficiency, as the discussion of distributed strategies demonstrates.

Within a node, GPUs may communicate through high-bandwidth links such as NVLink²⁷ (up to 900 GB/s aggregate bidirectional bandwidth for the represented H100 configuration). Effective collective bandwidth depends on topology and contention and is lower than a raw aggregate link specification. Data parallelism communicates gradients, while layerwise model parallelism communicates activations at partition boundaries. Choosing a strategy requires comparing its compute, memory, and communication on the actual topology.

8.7.2 Scaling beyond a single node

When single-node multi-GPU training remains insufficient, distributed training extends across machines. This introduces a slower or more contended communication tier and makes host, switch, routing, and fault behavior part of training performance. The bandwidth gap matters because synchronous data-parallel updates turn gradient reduction into a recurring data-movement problem.

🔍 Systems Perspective 8.8: The physics of synchronization

Recall the energy-movement invariant from chapter 4: data movement can cost substantially more energy and time than nearby arithmetic. In distributed training, this asymmetry contributes to the communication tax.

Synchronizing gradients across GPUs moves megabytes over a network or PCIe every few milliseconds. Uncovered communication lowers utilization, while communication energy raises the run’s energy budget; these are distinct costs. Mixed precision (section 8.6.3) and gradient

²⁷ | **NVLink:** NVIDIA’s high-bandwidth GPU interconnect. Its topology and generation determine the bandwidth available between a particular pair of GPUs and to collective operations. A parallelism strategy helps only when its communication and synchronization cost is smaller than the useful time it saves.

compression reduce bytes through narrower formats, sparsity, quantization, or encoding, addressing both pressures.

Within a node, fast links can make it possible to overlap part of gradient synchronization with backward computation. Crossing node boundaries adds another fabric whose effective bandwidth and latency depend on topology and collective implementation. In synchronous data parallelism, AllReduce²⁸ aggregates gradients across workers; at sufficient scale, its uncovered communication time can dominate the step.

A short calculation shows how quickly the inter-node fabric becomes the binding constraint.

Napkin Math 8.10: The network wall

Problem: A team scales training to 8 GPUs, one per node, connected by a 100 Gbps fabric. Is the network the bottleneck?

Mechanism: For a 7B-parameter model with FP16 gradients:

1. **Gradient size (FP16):** 14 GB per step.
2. **AllReduce cost:** Ring AllReduce passes gradient shards around a logical ring, so each worker ends up sending nearly two full copies of the gradient: $2(N_{\text{GPU}} - 1)/N_{\text{GPU}}$, which equals $1.75\times$ for 8 GPUs. That gives $1.75 \times 14 \text{ GB} = 24.5 \text{ GB}$ total.
3. **Bandwidth-only lower bound:** At 12.5 GB/s per worker on the 100 Gbps fabric, synchronization latency is $T_{\text{comm}} = D_{\text{comm}}/\text{BW}_{\text{net}} = 24.5 \text{ GB} / 12.5 \text{ GB/s} = 1.96 \text{ s}$.
4. **Compute comparison:** If the nonoverlapped forward and backward work takes $T_{\text{compute}} \approx 1 \text{ s}$, communication latency $T_{\text{comm}} \approx 1.96 \text{ s}$ dominates.

Systems lesson: The network becomes a wall when $T_{\text{comm}} > T_{\text{compute}}$, causing accelerator efficiency $\eta_{\text{hw}} = \frac{T_{\text{compute}}}{T_{\text{compute}} + T_{\text{comm}}}$ to plummet to about 33.8 percent. Mitigations include gradient compression (reducing D_{comm}), ring AllReduce primitives (Sergeev and Balso 2018), and hierarchical topologies that confine high-frequency transfers to fast intra-node links (NVLink at 900 GB/s vs 12.5 GB/s inter-node).

Once training crosses the node boundary, the strategy must match communication frequency to available bandwidth. Pipeline parallelism addresses idle time in layerwise partitioning through micro-batching. While one stage processes a later micro-batch, the next stage can process an earlier one. The schedule determines how many micro-batch activations each stage must retain until their backward computations; this can increase memory, but not every device stores the full model-wide activation footprint for every in-flight micro-batch. GPipe²⁹ and PipeDream developed influential schedules for this trade-off (Huang et al. 2019; Narayanan et al. 2019).

Tensor parallelism takes a finer-grained approach: rather than assigning whole layers to devices, it splits individual operations across devices. (Sequence parallelism extends this principle by sharding activation tensors along the sequence dimension across devices during long-sequence transformer training.) Consider a transformer's feed-forward layer with a large matrix multiplication $\mathbf{Y} = \mathbf{XW}$. Tensor parallelism splits the weight matrix $\mathbf{W}$ column-wise across GPUs, so each accelerator computes a portion of the output. The results are then gathered to form the complete output. This strategy is particularly effective for the massive attention and feed-forward layers in large transformers, where a single operation may involve matrices too large for one GPU's memory. Megatron-LM demonstrated that tensor parallelism enables training models with billions of parameters by distributing individual attention heads and feed-forward blocks across devices (Shoeybi et al. 2019).

Hybrid strategies combine these approaches because each has different scaling characteristics. Table 8.20 gives a topology-dependent heuristic: place the most frequent communication on the fastest links and less frequent synchronization on slower tiers.

The placement rule is not arbitrary; it maps communication frequency to the fastest available bandwidth tier.

²⁸ | **AllReduce:** A collective communication primitive that sums data across all devices and distributes the result back to each device. Ring AllReduce is one common implementation: devices pass gradient shards around a logical ring so each participant sends and receives bounded chunks rather than a full copy of every gradient tensor at once. This is the communication primitive that makes large data-parallel training practical, but its detailed cost model belongs with distributed training.

²⁹ | **GPipe (2019):** GPipe implements the described micro-batching strategy, staggering execution to fill the idle “bubbles” in naive model parallelism and thus boost utilization (Huang et al. 2019). The key trade-off it manages is increased memory for storing the activations of multiple in-flight microbatches. It preserves large-batch training dynamics by accumulating gradients before the weight update; the paper reports almost linear Transformer throughput scaling when the micro-batch count is at least four times the number of partitions.

Table 8.20: Hybrid Parallelism Placement: A common heuristic places tensor parallelism within nodes, pipeline parallelism within racks, and data parallelism across racks. Actual placement depends on the model and topology.

Strategy	Typical placement	Communication pattern
Tensor parallelism	Within a node	Frequent communication that exploits NVLink's high bandwidth
Pipeline parallelism	Across nodes within a rack	Moderate communication at layer boundaries
Data parallelism	Across racks	Gradient synchronization once per iteration

Data-parallel systems also rely on collective operations such as AllReduce to combine gradients across devices. The implementation details of collective algorithms, parameter-server and peer-to-peer communication patterns, fault tolerance mechanisms, and scaling-efficiency analysis for training runs spanning thousands of GPUs constitute a specialized domain that builds on the foundations established here.

8.7.3 The evolution of training infrastructure

The parallelism strategies in section 8.7.1 and section 8.7.2 are one endpoint of a longer infrastructure progression. Training systems took this form because computing infrastructure evolved through four distinct eras, each shaped by dominant workloads. Figure 8.19 situates those eras on a timeline, while table 8.21 compares their workloads, memory patterns, and processing models.

Neural-network training combines requirements from multiple predecessors. Like HPC, it uses substantial floating-point throughput; like warehouse-scale computing, large runs require coordination and fault handling across machines. Many current large-scale jobs use synchronous collectives each optimizer step, but asynchronous parameter servers, local-update methods, and other synchronization schemes also exist. The recurring high-volume communication of common synchronous training recipes helped motivate accelerator clusters with high-bandwidth interconnects and software stacks optimized for collective communication.




Figure 8.19: Computing System Evolution: Four overlapping eras place representative system milestones against time. The overlap matters: AI hypercomputing builds on, rather than replaces, the parallel numerical methods of HPC and the distributed coordination of warehouse-scale computing.

This architectural progression explains how training systems drew from earlier computing models. As table 8.21 shows, HPC systems provided the foundation for parallel numerical computation, while warehouse-scale systems demonstrated distributed processing at scale. Modern neural networks combine intensive parameter updates, complex memory access patterns, and coordinated distributed computation, prompting architectures that combine elements of both traditions.

In practice, configuring a multi-GPU training job entails choosing among parallelism strategies that evolved to address these distinct computational patterns. Understanding these strategies, their trade-offs, their communication costs, and their failure modes enables informed decisions about when additional hardware will help and when it will merely add complexity.

Table 8.21: Computing Era Characteristics: Each computing era optimized for different workload patterns, and modern training systems inherit requirements from multiple predecessors. AI hypercomputing uniquely combines HPC’s parallel numerical computation with warehouse-scale distributed processing, while adding specialized support for the gradient-based optimization and massive parameter state management central to neural network training.

Era	Primary Workload	Memory Patterns	Processing Model
Mainframe	Sequential batch processing	Simple memory hierarchy	Single instruction stream
HPC	Scientific simulation	Regular array access	Synchronized parallel
Warehouse-scale	Internet services	Sparse, irregular access	Independent parallel tasks
AI Hypercomputing	Neural network training	Parameter-heavy, mixed access	Hybrid parallel, distributed

8.7.4 When to scale: The physical ceiling

With this vocabulary of parallelism strategies (data, model, pipeline, tensor, and hybrid), knowing *how* to scale is different from knowing *when* to scale. Distributed training introduces substantial complexity. Before accepting it, practitioners should evaluate the single-machine techniques relevant to the measured constraint:

1. **Mixed-precision training:** Use supported lower-precision paths (section 8.6.3) when validation confirms acceptable numerical behavior.
2. **Gradient accumulation:** Use accumulation (section 8.6.5) to simulate larger batch sizes.
3. **Activation checkpointing:** Implement checkpointing (section 8.6.5.1) to trade compute for memory.
4. **Data pipeline optimization:** Optimize data pipelines (section 8.6.2) when input work lies on the critical path.

Table 8.22 turns this principle into a lookup keyed on measured resource limits rather than universal parameter-count thresholds.

Table 8.22: Scaling Decision Guidelines: Choose the least distributed configuration that satisfies measured memory, throughput, and schedule requirements. Model architecture, training state, input pipeline, hardware, and topology all affect the boundary.

Observed constraint	Candidate approach	Rationale
Training state fits one device and time is acceptable	Single GPU	Minimizes communication and operational complexity
State or throughput exceeds one device but fits one node	Single multi-GPU node	Uses fast local interconnects before crossing the network
Required state or schedule exceeds one node	Multi-node cluster	Adds aggregate memory or compute at a communication cost
Input service rate is below accelerator demand	Parallel or distributed input pipeline	Adds input throughput where profiling shows it is needed

Only when profiling reveals persistent bottlenecks despite these optimizations should multi-device approaches be considered. Every hardware device has a physical ceiling: a workload requiring 10^{24} FLOPs cannot be completed on a single accelerator within a practical training schedule, no matter how carefully that accelerator is tuned. The transition to multi-device training becomes necessary when one of three hard limits is reached:

- **Memory exhaustion:** The model weights, gradients, optimizer states, and activations exceed the VRAM of a single GPU. A 70-billion-parameter model requires approximately 140 GB in FP16 for weights alone, exceeding the 80 GB configurations of A100 and H100 SXM accelerators and leaving little or no headroom on higher-capacity variants like the H100 NVL (94 GB) and H200 (141 GB) once training state is included.
- **Training wall-clock time:** The estimated time on one device exceeds the project’s actual deadline. At 10^{15} FLOP/s sustained, a workload requiring 10^{24} FLOPs would take about 32 years.

- **Input throughput:** The input pipeline cannot sustain the aggregate accelerator consumption rate. Dataset size alone does not determine this limit; reuse, compression, locality, transformation cost, and storage bandwidth matter.

These three limits are technical, but scaling carries a cost beyond complexity and capacity. Adding devices amplifies the energy consumption and environmental impact of training, so efficiency optimization becomes an environmental concern as much as a performance one.



Napkin Math 8.11: The carbon footprint of training

Context: Training large models requires substantial energy and compute. The physical ceiling appears in the energy corollary to the iron law ($E_{\text{total}} = P_{\text{fleet}} \cdot T_{\text{train}}$), quantifying environmental scale:

1. Workload: Training a 7 billion-parameter model for 1 trillion tokens.
2. Compute: $O \approx 4.2 \times 10^{22}$.
3. Efficiency: 156 TFLOP/s sustained on A100 (400 W thermal design power).
4. **Execution time:** $T_{\text{train}} \approx 3.04$ days on 1,024 GPUs.
5. **Energy footprint:** $E_{\text{total}} = P_{\text{cluster}} \cdot T_{\text{train}} = (N_{\text{GPU}} \cdot P_{\text{GPU}} + N_{\text{host}} \cdot P_{\text{host}}) \cdot T_{\text{train}}$ evaluates to $(1,024 \text{ GPUs} \times 400 \text{ W} + 128 \text{ hosts} \times 200 \text{ W}) \times 73 \text{ hours} \approx 31,784.2 \text{ kWh}$.

Impact: The energy consumed equals 35.3 months of average household electricity use.

Systems lesson: Doubling hardware throughput η_{hw} at constant cluster power P_{cluster} halves execution time T_{train} and total energy E_{total} . Under average grid carbon intensity, doubling efficiency saves $\sim 15,892.1 \text{ kWh}$ of electricity and avoids 6.8 t of emissions. Environmental audits must measure $E_{\text{total}} = \int P(t)dt$ rather than relying on GPU utilization metrics alone.

Scaling changes the binding constraint rather than removing it: data parallelism buys throughput, model parallelism buys memory capacity, pipeline parallelism buys utilization, and tensor parallelism buys layer-scale feasibility, but each adds communication and coordination cost. The implementation details of multi-node distributed training, including collective communication primitives, fault tolerance mechanisms, and elastic scheduling, build directly on the single-machine principles covered throughout this chapter and are treated in depth in advanced distributed systems texts.



Checkpoint 8.3: Scaling decisions

Scaling trades compute bottlenecks for communication bottlenecks.

- ☐ **Hard limits:** can you identify which limit is binding for a given training run: memory capacity, wall-clock time, or input-service rate?
- ☐ **Single-node first:** can you explain why mixed precision, accumulation, checkpointing, and prefetching should be exhausted before adding devices?
- ☐ **Data vs. model parallelism:** given a model that does not fit on one GPU, can you justify why data parallelism fails and what model parallelism must partition?
- ☐ **Communication tax:** Can you explain why scaling efficiency degrades with GPU count, and what quantity (synchronization fraction) controls the ceiling?
- ☐ **Scaling ceiling:** if synchronization that does not parallelize is 10 percent of each step, estimate the speedup ceiling as the GPU count grows large, and the rough point where doubling devices stops being worth it.

8.8 Fallacies and Pitfalls

The progression from single-GPU optimization through multi-device parallelism shows that each technique introduces trade-offs and new constraints. The approach developed throughout this

chapter quantifies costs through the iron law, diagnoses bottlenecks through profiling, applies targeted optimizations, and scales only when necessary. The following fallacies and pitfalls capture common errors that waste compute resources, delay research progress, or cause production training failures.

Fallacy: *Larger models always yield better performance.*

Teams sometimes treat scale as a monotonic lever: if a 7B-parameter model works well, a 20B-parameter model must work better. A 20B model requires approximately 320 GB of nonactivation training state in the stated **mixed-precision Adam** recipe (40 GB FP16 parameters + 40 GB FP16 gradients + 80 GB FP32 master weights + 160 GB Adam states). Whether it outperforms a 7B model depends on data quantity and quality, training compute, optimization, and evaluation. No universal example-count threshold or accuracy penalty follows from parameter count alone.

Pitfall: *Assuming distributed training automatically accelerates development.*

More accelerators do not guarantee faster training because communication, smaller local batches, input limits, and imbalance can consume the gain. In an illustrative 8 GPUs scenario, synchronization occupies 30–50 percent of step time and speedup reaches only 4–6× rather than the ideal 8×. The values are hypothetical; the decision requires measured end-to-end time and scaling efficiency.

Fallacy: *Hyperparameters transfer directly from small-scale experiments to large-scale training.*

A learning rate that works at batch size 512 may not transfer to batch size 4,096. The linear scaling rule (Goyal et al. 2017) would increase the rate from 0.1 to 0.8 for this eightfold change, often with warmup, but it is an empirical recipe rather than a requirement or convergence guarantee. Large-scale runs must validate the learning rate, schedule, optimizer state, and effective batch together.

Pitfall: *Treating mixed precision training as a simple toggle without validation.*

The chapter’s illustrative V100 inputs imply a 2.4× speedup, but they are not a benchmark. FP16 commonly requires loss scaling (Micikevicius et al. 2017), while BF16 usually does not (Wang and Kanwar 2019; Kalamkar et al. 2019) (see section 8.6.3). Either recipe can change numerical behavior, so throughput and convergence must be validated on the target workload.

Fallacy: *Memory and computation can be optimized independently.*

Memory and compute are coupled: in this illustrative profile, accelerator utilization drops from 90 percent at batch 256 to 60–70 percent at batch 16. Gradient accumulation (effective batch 512, physical batch 64) trades 5 percent efficiency for 8× memory reduction. Tuning these parameters independently extends training time by 20–40 percent (see section 8.6.5).

Pitfall: *Budgeting training memory as if it only contains model weights.*

Sizing training memory from weights alone omits gradients, optimizer state, master weights when used, activations, temporary workspaces, and allocator overhead. In the chapter’s mixed-precision Adam recipe, FP16 weights (2 bytes), FP16 gradients (2), FP32 master weights (4), and two FP32 moments (8) total 16 bytes per parameter before activations. Other optimizers and sharding strategies change this budget. Activation memory then depends on batch size, sequence length, architecture, kernels, and checkpointing; no universal multiple of the weight footprint applies.

Fallacy: *The accelerator is always the training bottleneck.*

Data loading often creates idle time, yet teams optimize computation first. In this illustrative profile, prefetching with pipeline overlap reduces wall-clock time by 47.6 percent (105 min to 55 min) by overlapping data loading with computation (see section 8.6.2). Profile before assuming the GPU is the bottleneck.

Pitfall: *Optimizing kernels before profiling the input pipeline.*

Kernel-level tuning is attractive because GPU utilization is easy to inspect, but a starved accelerator can look like an inefficient accelerator. Before changing precision, rewriting kernels, or adding GPUs, measure dataloader throughput, host preprocessing time, storage wait, and host-to-device transfer overlap. If the accelerator is waiting for batches, the fastest optimization is to feed it reliably rather than make its kernels marginally faster.

8.9 Summary

Training combines algorithms, memory management, and hardware acceleration to transform data into model parameters. At scale, forward and backward propagation coordinate matrix operations, allocations, and gradients under hardware and performance constraints.

Prefetching, mixed precision, FlashAttention, gradient accumulation, and checkpointing address different throughput, memory, and numerical constraints. Match the technique to the binding constraint; when one machine is exhausted, data, model, pipeline, and tensor parallelism trade greater scale for greater complexity.

This co-design shapes large-scale training: matrix patterns drove GPU Tensor Cores, frameworks exposed them through mixed-precision APIs, and techniques such as FP16 training influenced later hardware. FLOP and memory accounting then support optimizer comparison, cost estimation, and decisions about whether more hardware will help or merely move the bottleneck.



Key Takeaways: Why training costs millions

- **Training cost is an iron-law budget:** $T_{\text{train}} = \frac{O}{R_{\text{peak}} \times \eta_{\text{hw}}}$ makes every optimization accountable: reduce work, raise effective throughput, or improve utilization. A change that misses the dominant term only moves cost around the training loop.
- **Memory determines whether an optimizer fits:** Standard mixed-precision Adam uses $8\times$ the FP16 inference-weight memory in the stated representation before activations. Optimizer state and batch-dependent activations—not weights alone—determine whether training fits.
- **Profiles choose the remedy:** Profile, diagnose, fix, and re-profile. Compute-bound jobs need better arithmetic or algorithms; memory-bound jobs need less state or traffic; data- and communication-bound jobs need pipeline or parallelism changes.
- **Precision and IO-aware kernels shift bottlenecks:** FP16 with FP32 accumulation can improve throughput and memory use, while FlashAttention avoids materializing the full $S \times S$ matrix in HBM; realized gains depend on workload and hardware.
- **Checkpointing buys memory with recompute:** Storing fewer activations and recomputing them during backpropagation cuts activation memory $3\text{--}4\times$ in the walkthrough, from 35.9 GB to 9 GB at batch 4. Use it when memory, not compute, binds.
- **Composed optimizations can postpone scale-out:** Mixed precision and checkpointing turn 95.7 GB into a modeled 33 GB for batch-4 GPT-2, before temporary workspaces. Use relevant local levers before accepting distributed communication and operational overhead.

The iron law classifies a slow job as compute-bound (improve arithmetic or algorithms), memory-bound (reduce state or traffic), data-bound (repair the input pipeline), or communication-bound (adjust accumulation or parallelism). Increase batch only when hardware is underfilled. This discipline prevents treating symptoms with more hardware and keeps iteration practical as cost grows.

Training is where the iron law becomes a daily instrument. Mixed precision changes effective throughput and memory traffic; checkpointing exchanges memory for recomputation; scaling adds communication. The engineering task is to identify the dominant term and apply the least costly intervention that changes it.



What's Next: From training to data selection

Training produces learned parameters by spending compute on every example; the chapter's optimized run still requires 33 GB of memory. Once the training loop is efficient, the next question is whether every example earns that cost. Repeated or low-information samples

consume the same forward pass, backward pass, optimizer update, and communication as examples that add coverage or correct a weakness.

Chapter 9 therefore turns upstream from execution to workload composition. It examines how redundancy, representativeness, uncertainty, and class coverage determine which examples should enter a finite training budget. A useful subset must reduce work without erasing rare cases or shifting the effective training distribution. When that balance holds, a carefully selected subset can approach full-dataset accuracy while lowering the cost of every epoch and making additional experiments affordable. The optimization target is no longer only how efficiently the system processes an example, but which examples the system should process at all.

III

OPTIMIZATION AND ACCELERATION

Optimization Principles

A working model is rarely an efficient one. Part II established how to construct ML systems that respect physical constraints; Part III addresses how to meet real-world demands on time, memory, and energy. Every optimization involves navigating a frontier: improving one metric (accuracy, latency, energy) while managing the cost to others. Every optimization in Part III enacts algorithm-machine co-design by trading mathematical structure for physical feasibility. The principles here define the physics of efficiency—the laws that determine *why* some models are fast and affordable while others are slow and prohibitively expensive.



Principle 5: The Pareto frontier

Invariant: After objective directions are normalized, a feasible point lies on the Pareto frontier exactly when no other point is at least as good on every objective and strictly better on one.

- **Quantization** trades numerical precision for reduced memory footprint.
- **Pruning** trades model capacity for smaller representations and can improve speed when the resulting sparsity or removed structures are supported by the target hardware.
- **Distillation** trades training compute for inference efficiency.

Implication: Systems engineers navigate this frontier to find the operating point appropriate to a specific deployment environment. There is no universal optimum.

Navigating the Pareto frontier requires knowing *which* resource to optimize. Before selecting a technique, engineers must diagnose whether the workload is limited by computation, memory movement, or dispatch latency. Arithmetic intensity supplies the first test.



Principle 6: Arithmetic intensity law

Invariant: Attainable throughput (R) is no greater than the minimum of peak compute (R_{peak}) and DRAM bandwidth (BW) times operational intensity at the cache–DRAM boundary (I) (Williams et al. 2009):

$$R \leq \min(R_{\text{peak}}, I \times \text{BW})$$

Operational intensity and DRAM bandwidth must both be measured at the cache–DRAM boundary; this ideal bound excludes launch, synchronization, and other fixed overheads.

Implication: Adding peak compute does not raise the ideal roofline bound for a bandwidth-bound kernel. Engineers must identify whether the bottleneck is compute, bandwidth, or another system term before selecting an optimization.

Many important ML kernels fall on the memory-bound side of the ridge point, especially low-reuse operations such as embedding lookup, normalization, softmax, depthwise convolution, and small-batch inference paths. Dense matrix multiplications and convolutions can instead be compute-bound when batching and hardware utilization are high. This split is a consequence of the memory wall: processor speed has historically outpaced memory bandwidth, and the cumulative gap has widened over three decades (Wulf and McKee 1995; Gholami et al. 2024). Neural networks, with

their massive weight tensors and uneven temporal locality, are especially vulnerable. The arithmetic intensity law diagnoses *where* a workload sits relative to this wall. The cost of moving data explains *why* the wall is so punishing. Table 8.23 maps each bottleneck type to representative optimizations that address it and to changes that leave the dominant term untouched.

Table 8.23: The Bottleneck Diagnostic: Before optimizing, identify which iron law term dominates. Optimizing the wrong term leaves the dominant bottleneck unchanged and usually yields little or no end-to-end improvement.

If the workload is...	Dominant Term	Optimization That Works	Optimization That is Wasted
Memory-Bound	D_{vol}/BW	Quantization, pruning, batching	Faster accelerator (more FLOP/s)
Compute-Bound	$O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$	Better kernels, Tensor Cores, faster accelerator	More memory bandwidth
Latency-Bound	L_{lat}	Kernel fusion, async dispatch, bounded microbatching under the latency SLO	More FLOP/s or bandwidth alone

Performance is only one cost of memory movement; energy exposes a second constraint.



Principle 7: The energy-movement invariant

Invariant: Total workload energy depends jointly on per-event energy and event counts. In the cited 45 nm reference, moving a 32-bit value from DRAM can cost roughly 100–1,000× more energy per event than a 32-bit floating-point operation, depending on operation type (Horowitz 2014). Here, $e_{\text{DRAM},32\text{-bit}}$ denotes the energy of one 32-bit DRAM access, and $e_{\text{arithmetic},32\text{-bit}}$ denotes the energy of one 32-bit floating-point arithmetic event in the cited 45 nm technology.

$$e_{\text{DRAM},32\text{-bit}} \gg e_{\text{arithmetic},32\text{-bit}} \quad \text{for the cited reference}$$

Implication: When memory movement dominates workload energy, optimization strategies should prioritize **kernel fusion** (avoiding intermediate off-chip traffic) and **quantization** (reducing data size); when computation dominates, reducing operation counts and improving arithmetic kernels remain central.

Even with perfect data locality and optimal bottleneck targeting, a final constraint limits how much speedup any optimization can deliver.



Principle 8: Amdahl's Law


Invariant: The maximum speedup of a system is limited by the fraction of the workload that cannot be accelerated (Amdahl 1967). Here, f_{parallel} is the fraction of baseline execution time accelerated and S_{parallel} is that fraction's speedup; the fixed-work model assumes the remaining fraction is unchanged and the acceleration introduces no additional overhead.

$$\text{Speedup} = \frac{1}{(1 - f_{\text{parallel}}) + \frac{f_{\text{parallel}}}{S_{\text{parallel}}}}$$

Implication: If 95 percent of a model runs 100× faster on a GPU, the total system speedup is capped at ~16.8×. This explains *why* **data loading** and **preprocessing** often become the ultimate bottlenecks in highly optimized systems.

Part III applies these principles systematically through the D·A·M taxonomy—Data, Algorithm, Machine—asking first whether the work is necessary, then whether it can be simplified, and finally how to do it faster (see Appendix A for the full diagnostic framework). Benchmarking closes the loop by testing whether those theoretical savings survive end-to-end execution without violating quality constraints.

Data Selection




9.1	Data Selection Fundamentals
9.2	Static Pruning
9.3	Dynamic Selection
9.4	Self-Supervised Learning
9.5	Synthetic Data Generation
9.6	Decision Framework
9.7	Selection Engineering
9.8	Cost Modeling
9.9	Distributed Selection
9.10	Cross-Layer Interactions
9.11	Measurement Framework
9.12	Fallacies and Pitfalls
9.13	Summary

Purpose

How can a carefully selected subset retain most of the value of a much larger dataset?


Training pays to acquire, label, score, move, and process examples even though their learning value differs. Large datasets contain redundant, noisy, or target-misaligned samples, while scarce regions may lack the cases a model needs. Data selection turns that heterogeneity into a systems optimization. It can remove avoidable work before training, direct labels and training steps toward informative cases as learning proceeds, or create missing examples when scarcity is the binding constraint. It can also invest once in reusable representation learning when that cost can be amortized across downstream tasks. The objective is therefore not the smallest dataset but the lowest end-to-end cost of reaching target quality without losing rare cases, underrepresented groups, or deployment-relevant coverage. Every strategy adds its own overhead: examples must be scored, generated, shuffled, stored, or validated, and those costs can erase the savings they promise. That judgment must be made end to end, because gains in labeling or training can be erased by selection latency, storage access, or distributed coordination. A method earns its place only when the learning it preserves or creates exceeds the work required to apply it. This distinction matters across repeated experiments, where useful reductions compound, and across adaptation, where the value of an example changes with the model and target distribution. Data engineering established that data is the source code of ML systems; data selection asks which evidence is useful now and which is still missing. In D·A·M terms, it is data-algorithm co-design: shape the training signal so each unit of machine work buys more useful learning.





Learning Objectives

- Explain data selection as data-algorithm co-design that reduces total operations before training begins
- Calculate information-compute ratio to decide whether additional examples improve learning per FLOP
- Compare deduplication, coreset selection, and quality pruning for reducing redundant pretraining data
- Design curriculum, active learning, or synthetic-data strategies for changing data value during training
- Apply the selection inequality to test whether selection overhead beats full-dataset training cost
- Evaluate distributed selection pipelines against storage locality, consistency, and GPU utilization constraints
- Select data-selection investments using ROI, amortization, and compute-optimal frontier diagnostics



Compute supply can outrun high-quality data supply.

1 | Scaling laws: Jared Kaplan and colleagues at Johns Hopkins and OpenAI empirically demonstrated in 2020 that language model loss follows power-law relationships with model size, dataset size, and compute budget over the regimes they studied (Kaplan et al. 2020). Their fitted data exponent was approximately $\alpha = 0.095$ in $\mathcal{L} \propto D^{-\alpha}$. The value is not universal, but the diminishing-return shape makes it possible to reason about when selection becomes more cost-effective than collection.

2 | Data wall: Unlike compute (which scales with capital expenditure) or algorithms (which improve through research), the stock of high-quality human-generated text grows slowly; Epoch AI's updated projections estimated that, under the paper's modeled consumption assumptions, models could use datasets roughly equal in size to the stock of public human-generated text between 2026 and 2032 (Villalobos et al. 2022). The point for systems design is not a fixed calendar deadline; it is that data can become a supply constraint rather than merely an economic one. This constraint directly affects the total operations (O) term: when quality data becomes scarce, additional compute yields diminishing returns regardless of hardware throughput.

9.1 Data Selection Fundamentals

TRAINING pays the iron-law cost for every example it processes, but not every example returns learning signal worth that cost. Data selection gives a clean, well-engineered dataset a systems objective: keep the examples that contribute the most learning per unit of compute. Data engineering makes the dataset reliable through correct labels, consistent schemas, and governed records. Data selection optimizes the dataset's value by extracting maximum learning from minimum samples, directly shrinking the total operations (O) term in the iron law (principle 3). The distinction matters: quality asks whether data is correct, while value asks whether correct data is worth the compute spent processing it.

Increasing the amount of training data has long been a standard way to improve models. Scaling laws¹ (Kaplan et al. 2020; Hoffmann et al. 2022) quantify how model performance can improve with dataset size, and teams have responded by collecting and generating more examples. Accelerator fleets, however, can expand usable compute faster than the supply of novel, high-quality human-generated text and images. Much of the easily accessible public web has already been incorporated into large training corpora, and expert labeling capacity grows slowly. This asymmetry is the **data wall**,² and it shifts attention from collecting more data to extracting more value from existing data. The sections that follow develop the selection methods and cost models needed to make that choice.

Table 9.1 uses an illustrative growth scenario in which GPU compute rises 10× every 3 years while accessible high-quality data grows more slowly. These are scenario inputs, not measured universal rates.

Table 9.1: Scaling Asymmetry in ML Resources: In the illustrated scaling regime, compute grows faster than high-quality data supply, creating a compute-to-data imbalance that makes data selection valuable.

Resource	Growth Rate	Implication
GPU Compute	~10× / 3 years	Hardware throughput can rise quickly in a given era
Training Data (Web)	~2× / 5 years	High-quality web text is finite; much already scraped
Labeled Data	~1.5× / 5 years	Human annotation throughput is inherently bounded
Synthetic Data	Potentially large	Bounded by generator quality (models trained on model-generated data can degrade)

Scaling dataset size eventually encounters the finite stock of high-quality human speech and text. Follow the log-scale trajectory in figure 9.1, observing how foundation model training requirements approach the estimated upper bound of available public data.

The gap between what compute can process and what accessible, high-quality data can support is therefore a systems variable, not a fixed law. In the regime illustrated by table 9.1, compute supply grows faster than accessible high-quality data, so accelerator budgets can outrun the corpus and leave the system compute-rich but data-constrained. Maintaining model relevance requires repeated refreshes of the training corpus as the world changes, and intelligent data selection becomes critical when data quality rather than accelerator time is the binding constraint.




Figure 9.1: Dataset Growth Approaching Limits: Several foundation-model datasets have grown toward estimates of the total stock of high-quality public text. The projection should be read as a dated constraint scenario rather than a fixed forecast: data selection, synthetic generation, multimodal learning, licensing, and filtering policy all change the usable supply.

The compute-data asymmetry can invert the optimization priority. When *data* is abundant and *compute* is scarce, algorithmic efficiency can extract more accuracy from limited GPU cycles. When *compute* is abundant and *quality data* is scarce, data selection can instead extract more learning from each sample. Data selection operates upstream of model and machine optimizations. By pruning redundancy and selecting high-value samples, the workload is reduced before it enters the model or reaches the hardware, directly shrinking the total operations (O) term in the iron law. That is why the systems perspective treats selection as upstream workload reduction rather than a modeling heuristic. For teams whose accelerator budget exceeds their curated corpus, the bottleneck shifts from GPU access to the quality, legality, and diversity of the training data.

The engineering toolkit for intelligent data selection follows a deliberate optimization ordering: first ask whether a sample is worth processing, then ask how to process the remaining workload efficiently. Data selection puts the “largest return first” principle into practice by looking for avoidable work before simplifying or accelerating the work that remains. Static pruning removes low-value samples before a single gradient is computed. Dynamic selection adapts the data diet during training through curriculum learning and active learning. Synthetic generation creates high-value samples through augmentation, simulation, or teacher-generated examples when real data runs short.

Each stage can increase the information density of the data that reaches the model when measured quality gains justify its overhead, and together they form a complementary toolkit: pruning reduces *what* the pipeline contains, selection focuses *how* the pipeline uses it, and synthesis expands *what* the pipeline can access. Comparing these techniques requires a systems definition of data selection and a measurable account of its effectiveness.

9.1.1 Defining data selection

The three-stage pipeline needs a quantity that lets engineers compare samples before they spend accelerator time on them. A duplicated image, a mislabeled record, and a rare boundary case may all cost the same forward and backward pass, but they do not contribute the same learning signal. Data selection therefore starts by making sample value explicit relative to compute cost.

The information-compute ratio measures that value as learning signal gained per unit of training compute, formalized in section 9.1.3.



Definition 9.1: Data selection

Data selection is the process of maximizing the information-compute ratio (ICR) of a training dataset.

1. **Significance:** It identifies data worth retaining, prioritizing, or generating for the target objective, reducing the total operations (O) when redundant or noisy samples can be excluded without losing required coverage.
2. **Distinction:** Unlike data engineering, which focuses on the cleanliness and consistency of data, data selection focuses on the informativeness and diversity of the samples.
3. **Common pitfall:** A frequent misconception is that more data is always better. Additional low-quality data can yield less improvement than a much smaller, carefully selected set, so retained volume must be evaluated together with coverage and target quality.

To make this concrete, consider training an autoregressive language model in the GPT-2/Llama lighthouse family from section 6.1.1. The scenario uses a 70-billion-parameter model to make the compute-data gap visible at large scale.

The compute budget (10,000 H100 GPUs for 3 months) represents roughly \$86.4M at the chapter's cloud-training price anchor and can process about 73.2T tokens at 40 percent sustained model FLOPs utilization. Estimates of quality- and repetition-adjusted public human-generated text are on the order of 300T tokens, far larger than a single curated web corpus. The practical bottleneck is narrower: how much of that stock is accessible, legally usable, high quality, deduplicated, and useful for the target distribution. If a team has only a 5T filtered corpus ready for use, the compute budget can already process it roughly 14.6× over. At that point, the team faces three options:

- **Repeat epochs:** The team can train on the same data for multiple epochs, but returns can diminish as examples are reused; the useful epoch count is workload-dependent.
- **Lower quality thresholds:** The team can include more data, but lower-quality tokens can degrade model quality.
- **Invest in data selection:** The team can improve filtering, curriculum design, and synthetic augmentation to extract more learning from each token.

Under these assumptions, the decision criterion favors selection over buying more accelerator time.

The opportunity applies across model architectures, though the bottlenecks and safe reduction differ. Unlike our compute-bound ResNet-50 lighthouse, GPT-2/Llama models can be memory-bandwidth-bound during inference and compute-bound during training. When a validated selection method removes examples or tokens without adding compensating steps, fewer samples mean fewer training FLOPs. The appropriate framing is therefore systemic rather than purely statistical.

9.1.2 Systems perspective

The data wall establishes *why* data selection matters; the systems perspective reveals *how* to approach it effectively. The conventional *ML framing* focuses on achieving the same accuracy with fewer samples, centering on statistical sample complexity and generalization theory. While valid, that framing misses the larger picture.

A data-selection *systems framing* asks instead how to reduce the total cost of achieving target performance across the entire ML lifecycle. The shift moves attention from accuracy curves to resource consumption, as table 9.2 illustrates.

Table 9.2: ML vs. Systems Perspectives on Data Selection: The ML framing optimizes sample complexity; the systems framing optimizes total resource cost across the pipeline.

ML Framing	Systems Framing
“Fewer samples for same accuracy”	“Fewer FLOPs for same accuracy”
“Better generalization”	“Lower training cost (time, money, energy)”
“Sample complexity bounds”	“End-to-end resource efficiency”
“Learning theory”	“Cost engineering”

The systems framing reveals optimization opportunities invisible to the ML framing. To see *why*, consider *how* data selection interacts with the iron law introduced in section 1.7.



Systems Perspective 9.1: Data selection and the iron law

In the iron law of ML systems ($T = \frac{D_{\text{vol}}}{R_{\text{BW}}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + I_{\text{lat}}$), data selection reduces the number of examples or tokens that enter the workload. Later model-level optimizations reduce operations per forward/backward pass, and machine-level optimizations increase R_{peak} (peak throughput) and η_{hw} (utilization). Selection removes work before those downstream optimizations execute it.

This makes data selection multiplicatively valuable in the iron law: when all three optimization layers act on the same bottleneck, a 2× reduction in dataset size with 2× fewer operations per sample and 2× higher effective throughput yields 8× total cost reduction, not 6×.

Consider training cost reduction. With the epoch count, batch size, and per-sample work held fixed, a 50 percent reduction in dataset size halves the number of forward passes, backward passes, and gradient updates. For the \$100M scenario used here, this translates to \$50M in compute savings. Schedules that add epochs or steps reduce that saving.

Compute savings cascade through the entire infrastructure stack. Large datasets consume petabytes of storage and saturate network bandwidth during distributed training; deduplication and coreset selection reduce storage costs while eliminating I/O bottlenecks that can idle expensive GPU clusters. The savings extend to labeling economics: expert labeling can exceed compute costs, and active learning and semi-supervised methods can substantially reduce labeling budgets in favorable regimes. The environmental implications compound further: for compute-dominated training runs, reducing the number of examples reduces energy consumption in proportion to the work avoided, provided the selected subset preserves accuracy. Smaller curated datasets also enable faster iteration velocity. A team that can iterate in hours rather than days has a compounding advantage in model development.

The cascading benefits illustrate a broader point: the ML researcher usually frames the problem as sample complexity, while the systems engineer frames it as cost-per-accuracy-point across the entire pipeline, from data acquisition through deployment. The systems engineer’s toolkit for that problem includes techniques to minimize total cost, metrics to quantify efficiency gains, and architectural patterns to implement data selection at scale.

9.1.3 Information-compute ratio

The systems framing established in section 9.1.2 calls for a quantitative metric. Data selection creates a frontier between accuracy and cost: keeping every example maximizes coverage but wastes compute on redundancy, while pruning too aggressively saves compute but loses signal. We need a way to measure where a sample sits on that frontier. The metric is the information each sample contributes to the model’s learning per unit of computation. We formalize it as the **information-compute ratio**.

Figure 9.2 recasts the D·A·M taxonomy as an optimization map, with data selection playing the role of input optimization: reducing total workload before it enters the model or hardware. The model side asks how much math each example requires. The machine side asks how quickly the hardware can execute that math. The data side asks whether the example should be processed

at all. The three edges of the triangle capture the dominant bottlenecks: *compute bound* describes systems limited by arithmetic throughput, *I/O bound* describes systems limited by data movement, and *sample efficiency* describes systems limited by the information content of training data.




Figure 9.2: The D·A·M Taxonomy as an Optimization Map: Machine learning performance is organized around Algorithms (Model), Data Selection, and Machine (Hardware). The triangle's edges label the compute-bound, I/O-bound, and sample-efficiency relationships, while data selection reduces workload before model or hardware execution.

We can formalize this as the ICR, where I denotes information content:

$$\text{ICR} = \frac{\Delta I}{\Delta \text{FLOPs}}$$

A higher ICR means each FLOP of training buys more learning; pushing it up is the goal of every technique in this chapter. The numerator is not directly observable in production, so engineers estimate it through proxies: validation improvement per unit compute, area under the learning curve, loss reduction on held-out data, uncertainty or gradient-based sample scores, and coverage checks on deployment-relevant slices. Those proxies are imperfect, but they make the systems question measurable before the training run spends its full budget.


9.1.4 The ICR frontier: When data becomes a tax

The information-compute ratio is not constant; it follows a law of diminishing returns. The **ICR frontier** is the point where the marginal learning signal from additional data drops toward zero.

To illustrate diminishing returns, let $I(D)$ be the information content of a dataset of size D and assume $I(D) \propto \log D$. This is a simplified analytical model rather than a universal law. If compute cost scales linearly with the per-sample operation count O_{sample} , then $C(D) = O_{\text{sample}} \cdot D$, and the resulting ICR follows equation 9.1:

$$\text{ICR}(D) = \frac{\frac{d}{dD} I(D)}{\frac{d}{dD} C(D)} \approx \frac{1/D}{O_{\text{sample}}} = \frac{1}{O_{\text{sample}} \cdot D} \quad (9.1)$$

Under this illustrative model, the $1/(O_{\text{sample}} \cdot D)$ decay creates the data wall. Beyond a workload-dependent frontier, additional data may provide little learning while still adding compute. In this regime, data becomes a **data tax** that inflates the O term of the iron law without a commensurate improvement in the accuracy numerator of the RoC (return on compute, see section 1.7.1). The knee is a practical trade-off between target performance and marginal cost, not the mathematical maximum of ICR in this model.



Past the frontier, data becomes a tax: compute climbs, learning stalls.

Data selection turns the total operations (O) term in the iron law from a fixed constant into a variable. A $2\times$ improvement in measured ICR halves the FLOPs required to reach the same measured gain. It produces the same ideal time reduction as a $2\times$ throughput increase only when the workload is compute-bound and the saved work lies on the critical path. ICR focuses specifically on compute; the cost-modeling framework in section 9.8 extends the same reasoning to acquisition, labeling, and storage costs.

A random batch of raw data can have low ICR when it contains redundant, noisy, or already-mastered examples. High-efficiency data pipelines (figure 9.3) can improve ICR through three stages: static pruning before training, dynamic selection during training, and synthetic generation on demand. To illustrate, consider computing ICR on a concrete coreset selection task: a deliberately selected subset intended to preserve the full dataset’s learning signal. Section 9.2.2 defines the EL2N and GraNd scoring methods used to build such subsets, and section 9.11 provides the complete measurement framework for evaluating these efficiency gains, including a compute-optimal frontier diagnostic that can suggest whether training is data-limited or compute-limited.




Figure 9.3: The Data Selection Pipeline: A structured approach to increasing data value. Raw data is first pruned (Static Pruning), then selected during training (Dynamic Selection), and finally synthesized (Synthetic Generation). Each stage can increase the Information-Compute Ratio (ICR) when its quality gain justifies its compute cost.



Checkpoint 9.1: Data selection efficiency

The goal of data selection is to maximize the ICR.

Metric checks:

- ☐ **ICR application:** given two training runs with identical accuracy gains but different compute budgets, can you determine which had higher ICR?
- ☐ **Data efficiency:** under a fixed evaluation protocol, can you explain when a 50 percent smaller dataset with $2\times$ higher ICR yields the same measured gain for half the training FLOPs?

Pipeline check:

- ☐ **The three stages:** can you map static pruning, dynamic selection, and synthetic generation to the training lifecycle?

The practical question is how large the efficiency gap could become on a workload where dataset size, model cost, and selection strategy interact with concrete FLOP budgets. The following ImageNet-scale ResNet-50 scenario uses assumed accuracy gains to illustrate the calculation; it is not a reported EL2N result.



Napkin Math 9.1: Computing ICR: Coresets

Problem: Training the ResNet-50 lighthouse model from section 6.1.1 on ImageNet for one epoch costs 3.15×10^{16} FLOPs. Under the scenario’s assumed accuracy gains, how would retaining 50 percent of the data change the information-compute ratio? The calculation holds the per-sample work fixed and does not include selection overhead.

Setup:

- Dataset: ImageNet (1.28M)

- Model: ResNet-50 lighthouse (~8.2 GFLOP per forward pass, roughly 24.6 GFLOP for a forward plus backward training step, depending on implementation)
- One epoch: $1.28\text{M} \times 24.6 \text{ GFLOP} = 3.15 \times 10^{16} \text{ FLOPs}$
- Assumed accuracy improvement per epoch (early training): 5 percentage points

Random selection (baseline):

- Process all 1.28M samples uniformly
- Accuracy gain: 5 percentage points
- $\text{ICR}_{\text{random}} = 5 \text{ percentage points} / (3.15 \times 10^{16} \text{ FLOPs}) = 1.6 \times 10^{-16} \text{ per FLOP}$

EL2N coreset (Error L2-Norm, a training-dynamics score developed in section 9.2.2; 50 percent of data):

- Process 640.6K selected samples
- Assumed accuracy gain: 4.5 percentage points (90 percent of the full-data gain)
- Compute: $640.6\text{K} \times 24.6 \text{ GFLOP} = 1.6 \times 10^{16} \text{ FLOPs}$
- $\text{ICR}_{\text{coreset}} = 4.5 \text{ percentage points} / (1.6 \times 10^{16} \text{ FLOPs}) = 2.9 \times 10^{-16} \text{ per FLOP}$

Systems insight: Under these assumptions, the selected subset achieves 1.8× higher ICR with a 0.5 percentage points smaller accuracy gain. A real decision requires measured gains, selection overhead, and coverage checks.

The three-stage optimization pipeline (static pruning, dynamic selection, and synthetic generation) provides concrete techniques for improving ICR. Static pruning, the first stage, removes examples before training begins; the safe reduction depends on the dataset, model, selection method, and target metric.

9.2 Static Pruning

The first stage of the pipeline acts entirely before training begins, removing low-value samples so that fewer of them ever reach the model. Static pruning and pretraining filtration can reduce total computation without modifying the training loop or model architecture, but their effect on final quality must be measured.

9.2.1 The case for smaller datasets

Data-selection studies show that some datasets can be reduced without lowering measured accuracy. More data can improve performance, but large datasets may also contain substantial redundancy. Empirical studies on coreset selection and data pruning have demonstrated this effect on several standard benchmarks.

On CIFAR-10, Paul et al. (2021) report that EL2N pruning removed half of the training data without reducing accuracy. On ImageNet-1K, Sorscher et al. (2022) report that a self-supervised prototype metric discarded 20 percent of the data without sacrificing performance. The pattern extends to language modeling, where web-scraped corpora like The Pile³ and C4⁴ contain enough duplicate and templated content to make deduplication a systems issue. In datasets studied by Lee et al. (2022), approximate near-duplicate removal affected 3.04 percent of C4 and 13.63 percent of RealNews, and deduplicated training reduced memorization while preserving or improving perplexity.

The reported gains are benchmark-specific. Pruning effectiveness depends on the dataset’s intrinsic redundancy, the selection algorithm, and the model architecture; validate the result on the specific task before deploying aggressive pruning in production.

Individual data points need not provide equal value for training. This heterogeneity follows in part from how classifiers learn decision boundaries. Many samples fall far from a boundary: a picture of a dog in good lighting may become easy once the model has learned its dominant features, while ambiguous and underrepresented cases can remain informative. Label quality also affects data requirements. In a simplified binary symmetric-noise model, a label-flip rate ρ reduces an

³ | **The Pile:** An approximately 886 GB English text corpus (reported as 825 gibibytes by the source) aggregating twenty-two sub-datasets, including PubMed, ArXiv, GitHub, Project Gutenberg, Common Crawl, Stack Exchange, Wikipedia, and USPTO patents (Gao et al. 2020). Its multi-source design makes it a useful example of data diversity: each source family has a different duplication, quality, and domain-coverage profile, so selection pipelines must preserve coverage while removing redundant text.

⁴ | **C4 (Colossal Clean Crawled Corpus):** C4 applies filtering to Common Crawl data, including language detection, deduplication of repeated three-sentence spans, and removal of pages that fail heuristic quality checks, producing approximately 750 GB of cleaned English text (Raffel et al. 2020). Filtering has its own compute cost and can remove useful material, so its ROI and quality effects must be measured rather than inferred from corpus size alone.

effective-sample-size factor to $(1 - 2\rho)^2$; at $\rho = 10\%$, matching the clean-data bound requires about $1/(1 - 2\rho)^2 \approx 1.56\times$ as many samples. The notebook below uses a separate toy learning-curve model to illustrate another **data quality multiplier**; neither multiplier is universal.

Napkin Math 9.2: The data quality multiplier

Problem: How much more noisy data does it take to reach the same target error as clean data?

Math: In this illustrative model, clean-data error scales as $\mathcal{O}(1/D)$, so $D_{\text{clean}} \propto 1/\epsilon$. Noisy-data error scales as $\mathcal{O}(1/\sqrt{D})$, so $D_{\text{noisy}} \propto 1/\epsilon^2$. For target error $\epsilon = 0.01$ (1 percent):

- $D_{\text{clean}} \approx 100$
- $D_{\text{noisy}} \approx 10,000$

Result: Under these assumptions, noisy data requires 100× more samples at the target error.

Systems insight: Here, cleaning data acts as a 100× compute accelerator; the real multiplier is workload-dependent.

9.2.2 Coreset selection algorithms

The practical question then becomes how to identify which samples to keep. **Coreset selection**⁵ turns the static pruning decision into a coverage problem: keep the smallest subset that preserves the statistical properties of the entire dataset.

The systems decision is where to spend the selection budget: on cheap coverage metrics that preserve distributional structure, or on costlier training-dynamics scores that better target the decision boundary. The decision also needs a guardrail before any score is applied: classes, demographic groups, time windows, and rare failure modes that matter at deployment require minimum representation, because the highest-average-ICR subset can still remove the examples that define production risk.

Geometry-based methods select samples that cover a chosen representation without requiring target-model training. The k -Center algorithm,⁶ a facility-location-style objective, selects samples that reduce the maximum distance from any point to its nearest selected center under the chosen embedding and metric.

Herding takes a different approach. Welling’s original procedure iteratively constructs representative pseudo-samples whose feature statistics approximate observed moments (Welling 2009). Related moment-matching objectives can guide subset selection from a fixed pool. These methods are computationally attractive because they operate on feature representations, but their scores are not inherently label-aware.

Training-dynamics methods use a proxy model to identify examples that may be important. GraNd (Gradient Normed) and **EL2N (error L2-norm)**⁷ score samples by gradient magnitude or prediction error early in training (Paul et al. 2021). High scores identify examples the proxy finds difficult, but they can also identify noise or unrepresentative samples and therefore need coverage and quality checks. The same work reports useful score transfer across several architectures, which can enable less expensive proxy-based selection. Forgetting Events⁸ tracks how often a sample is correctly classified and later misclassified during training (Toneva et al. 2019).

Geometry-based and training-dynamics methods optimize different proxies, so their quality ranking depends on the dataset, representation, target model, and scoring budget. The comparison should be read as a selection-budget table: a score is useful only if its target-quality gains justify its scoring cost.

Each algorithm in table 9.3 occupies a different point in the ICR framework’s compute-versus-information trade-off. Figure 9.4 illustrates one uncertainty-based strategy rather than every coreset method. Random sampling appears on the left, while the right panel concentrates the selection budget near the illustrated decision boundary. Geometry-based methods may instead prioritize coverage across the embedding space.

⁵ | **Coreset (core set):** In computational geometry, a small subset can approximate a specified geometric objective within a controlled error factor (Agarwal et al. 2005). Machine learning adapts this idea as a selection principle, but a guarantee for an embedding and distance objective does not by itself guarantee neural-network accuracy. The retained subset must still be validated on the target task.

⁶ | **k-center algorithm:** Its greedy strategy iteratively picks the point farthest from the existing centers under a chosen embedding and metric. Sener and Savarese use this core-set framing for convolutional-neural-network active learning (Sener and Savarese 2018). The resulting coverage bound applies to that geometric objective, not directly to downstream accuracy or rare-class preservation.

⁷ | **EL2N (error L2-norm) and GraNd (gradient normed):** These methods score samples based on their error or gradient norm early in training, identifying samples the model finds most difficult. The practicality of this approach relies on *transferability*, where scores from a small proxy model can guide data selection for a much larger target model. For instance, a proxy trained for five epochs can generate scores to curate a dataset for a full 90-epoch production training run (Paul et al. 2021).

⁸ | **Forgetting events:** This method identifies valuable examples by tracking when the model “forgets” them—transitioning from a correct to an incorrect classification during training. The central trade-off is the high cost of this analysis, which requires a full training run. However, the resulting importance scores transfer reliably from small proxy models to large target models (for example, ResNet-18 to ResNet-50), which is precisely what makes the “inexpensive proxy-based selection” strategy viable (Toneva et al. 2019).

Table 9.3: Coreset Selection Algorithm Comparison: D = dataset size, K = coreset size. Geometry-based and training-dynamics methods optimize different proxies, so their quality ranking depends on the dataset, representation, target model, and scoring budget.

Method	Compute Cost	Requires Training	Best For	Limitation
k-Center	$\mathcal{O}(D^2)$ or $\mathcal{O}(DK)$	No	Coverage, exploration	Ignores label information
Herdning	$\mathcal{O}(DK)$	No	Distribution matching	Assumes Gaussian-like
GraNd	$\mathcal{O}(\text{epochs} \times D)$	Yes (few epochs)	Decision boundaries	Requires proxy training
Forgetting	$\mathcal{O}(\text{full training})$	Yes (full)	Hard examples	Expensive to compute
EL2N	$\mathcal{O}(\text{epochs} \times D)$	Yes (few epochs)	Uncertainty sampling	Best with proxy model




Figure 9.4: Illustrative Uncertainty-Based Selection: Both panels select five samples, isolating allocation rather than budget. Random sampling distributes those choices across the feature space; the illustrative uncertainty strategy concentrates them in the shaded band around the current decision boundary.

EL2N with a small proxy model offers one practical balance between scoring cost and selection quality. The approach trains a lightweight model for a few epochs, computes EL2N scores, and selects a subset subject to coverage and quality checks. The proxy must provide rankings that transfer to the target model, and the investment pays off only when downstream training savings exceed selection overhead. A concrete scenario illustrates this workflow without asserting that its 10 percent subset preserves accuracy.



Example 9.1: Coreset selection in practice

Scenario: An engineering team aims to select a 100,000 coreset (10 percent) from 1 million training images for faster model iteration.

Diagnosis: Naive random subsampling drops rare classes and boundary examples. Training a lightweight proxy model for 5 epochs calculates EL2N error scores, retaining high-uncertainty samples near the decision boundary.

Systems lesson: Proxy-based coreset selection replaces some full-data passes with targeted scoring compute. It accelerates full-model training only when the proxy ranking transfers and coverage checks preserve rare classes and edge cases.

Listing 9.1 demonstrates how to compute EL2N scores and select a coreset using a lightweight proxy model. The mechanism spends a small amount of probe compute to identify high-uncertainty samples near the decision boundary, then trains the full model on the retained subset rather than on redundant easy examples. The example requires an unshuffled dataloader so each score position remains aligned with its dataset index.

Proxy scoring turns uncertainty into reusable selected indices: `compute_el2n_scores` measures which samples still confuse a briefly trained model, and `select_coreset` retains those high-information examples for the full training run.

Listing 9.1: EL2N-Based Coreset Selection: The `compute_el2n_scores` function trains a proxy model for a few epochs, then measures prediction error via L2 distance from one-hot labels. High scores identify samples the proxy finds difficult, including potentially noisy examples. The `select_coreset` function illustrates ranking by this score; production use also requires coverage, noise, and target-quality validation.

```
def compute_el2n_scores(model, dataloader, num_epochs=5):
    """Compute EL2N scores.

    Returns L2 norm of (prediction - one_hot_label).
    """
    # Train proxy model for a few epochs to get meaningful predictions
    train_proxy(model, dataloader, num_epochs)

    scores = []
    model.eval()
    for x, y in dataloader:
        logits = model(x)
        probs = softmax(logits, dim=1)
        # One-hot encode labels
        one_hot = zeros_like(probs).scatter_(1, y.unsqueeze(1), 1)
        # EL2N score = L2 distance from confident prediction
        el2n = (probs - one_hot).norm(dim=1) # High = uncertain
        scores.extend(el2n.tolist())
    return scores

def select_coreset(scores, dataset, fraction=0.1):
    """Select top-k highest-scoring (most uncertain) samples."""
    k = int(len(dataset) * fraction)
    # Sort by score descending (highest uncertainty first)
    indices = argsort(scores, descending=True)[:k]
    return Subset(dataset, indices)

# Illustrative 10% subset; validate accuracy and coverage before use
scores = compute_el2n_scores(proxy_model, full_loader)
coreset = select_coreset(scores, full_dataset, fraction=0.1)
train_full_model(model, coreset)
```

9.2.3 Data deduplication

While coreset selection identifies which samples to keep based on their informativeness, a complementary approach targets duplicate samples that add compute without adding learning signal. Deduplication can provide immediate efficiency gains, especially for exact and near-duplicates, and requires no model training. This makes it one of the most accessible optimizations in data selection, but near-duplicate thresholds must be validated so the pipeline does not remove useful distributional signal.

The simplest form of deduplication (introduced as a data engineering pipeline stage in section 4.6, and here elevated to an optimization lever) uses hash-based methods for exact matches. By computing a cryptographic hash (MD5 or SHA-256) for each sample and removing those with identical hashes, practitioners can eliminate byte-for-byte duplicates that inevitably accumulate in large web-scraped corpora. This process is computationally cheap, scaling linearly with dataset size, and can be parallelized trivially.

Near-duplicate detection addresses content that differs at the byte level while retaining substantial token- or character-shingle overlap. For text, MinHash⁹ with Locality-Sensitive Hashing¹⁰ (LSH) approximates Jaccard similarity¹¹ efficiently. It can detect lightly edited content with high overlap, but it does not establish semantic paraphrase equivalence.

For images, perceptual hashing produces signatures robust to minor transformations like resizing and compression, identifying visually similar candidates stored in different formats. Embedding-based similarity can retrieve semantic near-neighbor candidates by computing dense representations (CLIP¹² (Radford et al. 2021) for images, sentence transformers for text), though proximity does not itself prove duplication and the approach incurs higher computational overhead.

⁹ | **MinHash:** Invented by Broder (1997) to detect duplicate web pages for AltaVista, the algorithm creates compact signatures using random hash functions such that similar documents produce similar signatures with high probability. Each signature compresses a document to a fixed-size sketch, enabling pairwise similarity estimation in $\mathcal{O}(s)$ time for sketch size s .

¹⁰ | **Locality-sensitive hashing (LSH):** LSH hashes MinHash signatures into buckets so that similar signatures are more likely to collide. Candidate generation avoids an exhaustive all-pairs comparison, while runtime and memory depend on signature width, banding, bucket sizes, and the number of candidates that require verification.

¹¹ | **Jaccard similarity:** Defined as $|A \cap B| / |A \cup B|$, ranging from 0 (disjoint) to 1 (identical). For deduplication, the metric compares sets of shingles across documents of different lengths. The threshold must be calibrated for the shingle definition, corpus, and downstream quality because no universal cutoff separates duplicates from legitimately related documents.

¹² | **CLIP (Contrastive Language-Image Pretraining):** Pretrained on 400 million image-text pairs, CLIP maps semantically related images and text into a shared embedding space (Radford et al. 2021). Embeddings can retrieve candidates for review, but semantic proximity is not proof that two samples are duplicates. Generating embeddings also costs substantially more than computing perceptual hashes, with the ratio depending on model and hardware.

Deduplication can reduce repeated-content memorization and avoid spending FLOPs on exact or near duplicates. It can also alter the corpus’s intentional empirical weighting, so thresholds and downstream quality must be validated rather than assuming that every duplicate is harmful.

Deduplication also changes when retained examples index a stateful representation rather than only a training corpus.



Lighthouse 9.1: DLRM and embedding deduplication

Our DLRM lighthouse model from section 6.1.1 is memory capacity-bound, with embedding tables consuming terabytes of storage for billions of user/item IDs. Much of this capacity is wasted on *cold embeddings*, IDs that appear rarely in training data.

Data selection for DLRM can include interaction deduplication and embedding pruning for cold IDs. Removing repeated interactions reduces training examples but shrinks an embedding table only when IDs disappear or the allocation, hashing, or sharing policy changes.

9.2.4 Data pruning by quality

Deduplication removes redundant samples, but a third category of problematic data remains: samples that actively harm learning. Quality-based pruning eliminates samples that either contribute no meaningful signal or introduce contradictory information that confuses the optimization process.

Label error detection can be a high-impact form of quality pruning. Tools like Cleanlab identify samples where the assigned label may be incorrect based on model confidence patterns across training. A sample that the model consistently predicts as class A but is labeled class B may be a hard boundary case, a distribution shift, a model blind spot, or an annotation mistake. The score should trigger review rather than automatic deletion; confirmed errors can then be corrected or removed.

Outlier removal addresses a different pathology: samples far from any cluster center in feature space. Such points can indicate noise, annotation errors, or corruption, but they can also be the rare edge cases that define deployment risk. The key is distinguishing informative outliers from invalid samples, with conservative thresholds and slice-level review before removal.

Low-information filtering applies domain-specific heuristics to remove samples that lack sufficient signal for learning. For text corpora, this often means removing high-perplexity garbled text (perplexity is a language model’s measure of how surprising a text is, so high values flag incoherent strings) and, in some pipelines, low-perplexity boilerplate or repetitive text. For image datasets, filtering targets blurry, corrupted, or near-uniform samples that provide little visual information.

Together, these three static pruning techniques—coreset selection, deduplication, and quality filtering—can reduce work before training. With epoch count, batch size, and per-sample work fixed, a 50 percent dataset reduction means 50 percent fewer forward passes, backward passes, and gradient updates. Training schedules that add compensating epochs or steps reduce that saving.

Static pruning answers a question about *what* to keep, but it treats the answer as fixed. Once the pruned dataset is set, every epoch trains on the same subset. A sample’s usefulness, however, can change as the model learns: examples that challenge an undertrained model may become easy after sufficient gradient updates. Dynamic selection techniques address this limitation by adapting the training data at each stage based on what the model has already mastered.

9.3 Dynamic Selection

Early in training, a model benefits from broad, easy coverage that builds stable feature representations. Later, those same examples produce little new gradient signal, while harder samples near the decision boundary become more valuable for refinement. Dynamic selection exploits this changing information-compute ratio by adapting the data diet to the model’s current state.

9.3.1 Curriculum learning: Easy to hard

The first dynamic selection technique, **curriculum learning**¹³ (Bengio et al. 2009; Soviany et al. 2022), structures the order in which data is presented to the model. Instead of random shuffling, it starts

¹³ | **Curriculum learning:**

From Latin *currere* (“to run”), originally meaning “the course to be run”—a metaphor that maps directly to the technique: training data as a course run in deliberate order, easy stretches first. Curriculum learning can act as a continuation method for nonconvex optimization, with easier examples shaping the early optimization trajectory before harder examples are introduced. From a systems perspective, the ICR of a sample can vary during training, which motivates changing the mix of examples over time.

with simpler examples and gradually introduces more complex ones, mirroring how humans learn by mastering basics before advancing to harder material.

Curriculum learning can change the optimization trajectory by controlling which gradients dominate at each stage. Easy examples may provide stable early signals, while examples scored as hard can contain useful boundary information, label noise, or outliers. An easy-to-hard schedule can smooth training on some workloads, but its benefit depends on the difficulty score, pacing rule, model, and data; table 9.5 therefore treats the savings as illustrative rather than guaranteed.

Implementing a curriculum requires two components: a difficulty scorer that ranks samples, and a pacing function that controls how quickly hard samples are introduced. A common choice is linear pacing:

$$\text{samples}_{n_{\text{epoch}}} = \text{sort_by_difficulty}[: D \cdot \min(1, n_{\text{epoch}}/N_{\text{warmup}})]$$

where n_{epoch} is the current epoch, D is the full dataset size, and N_{warmup} is the number of warmup epochs before the full dataset becomes available. Early epochs train on the easiest $D \cdot (n_{\text{epoch}}/N_{\text{warmup}})$ fraction; after warmup, training proceeds on the full dataset.

The difficulty scorer is a systems choice because it trades probe-compute overhead against ordering quality, as table 9.4 shows. Loss and confidence scoring buy better ordering with extra inference, heuristics avoid compute but require domain knowledge, and self-paced scoring moves adaptation into the training loop itself.

Table 9.4: Difficulty Scoring Strategies for Curriculum Learning: The rows compare each scorer’s signal source with the system prerequisite or trade-off it introduces.

Strategy	Difficulty Score	Best For
Loss-Based	Loss from probe model (low = easy)	General-purpose; requires probe training
Confidence-Based	Teacher model confidence (high = easy)	When teacher available; distillation setups
Domain Heuristics	Sentence length, image complexity	No extra compute; domain knowledge required
Self-Paced	Current model’s loss (updated each epoch)	Adaptive; no probe needed

Table 9.5 uses hypothetical epoch counts to illustrate how curriculum savings would be computed across workloads. The values are not measurements from the cited curriculum-learning studies.

Table 9.5: Illustrative Curriculum Epoch Reductions: The table applies one target-accuracy rule to hypothetical baseline and curriculum epoch counts. It demonstrates the arithmetic but does not establish a ranking among datasets or methods.

Dataset	Model	Pacing Strategy	Epochs to Target Acc.	Epoch Reduction
CIFAR-10	ResNet-18	Linear warmup	115 vs. 150 baseline	23.3% fewer epochs
CIFAR-100	ResNet-32	Self-paced	180 vs. 220 baseline	18.2% fewer epochs
ImageNet	ResNet-50	Loss-based	80 vs. 90 baseline	11.1% fewer epochs
ImageNet	ResNet-50	Learned (noisy labels)	70 vs. 90 baseline	22.2% fewer epochs

The scenario illustrates that useful curriculum gains must be measured at a fixed target metric. The ordering is task-dependent: *anti-curriculum* presents hard examples first, while *self-paced learning* adjusts difficulty using the model’s current loss. Neither ordering is universally superior.

9.3.2 Active learning: Human-in-the-loop

Curriculum learning optimizes the order in which samples are presented but assumes all samples are already labeled. This assumption breaks down in specialized domains where labeling requires substantial expertise, time, or money. Rather than labeling everything upfront, active learning¹⁴ (Settles 2012; Ren et al. 2021) shifts the optimization target from choosing which labeled samples to train on to choosing which unlabeled samples are worth labeling.

Active learning transforms data acquisition from static collection into an iterative query process. Trace the circular flow path in figure 9.5, following how model uncertainty metrics select high-value unlabeled samples for expert annotation.

¹⁴ | **Active learning:** The “active” component is the learning algorithm itself selecting which unlabeled samples a human expert should label next. This reframes the problem from a computational optimization (training on given data) to a financial one: maximizing model improvement per dollar spent on expert labeling. Querying the most informative examples can substantially reduce labeling needs compared with random sampling, but the multiplier is task-, model-, and oracle-dependent.

The effectiveness of active learning depends critically on the query strategy used to select samples for annotation (Settles 2009, 2012; Ren et al. 2021). The simplest approach, uncertainty sampling, selects samples where the model is least confident, such as predictions near 0.5 probability for binary classification. It is inexpensive to score with one model but can overconcentrate queries near one ambiguous region. Query-by-committee extends this idea by training multiple models and selecting samples where they disagree most, capturing epistemic uncertainty that a single model might miss.

For practitioners willing to invest more compute, expected model change selects samples that would cause the largest gradient update if labeled. This approach provides a theoretically grounded but expensive alternative. Diversity sampling complements uncertainty-based methods by selecting samples dissimilar from currently labeled data, improving coverage rather than clustering every query around the same ambiguous region.

Active learning is particularly valuable in domains where labeling requires expertise. In medical imaging, for instance, an AI system diagnosing diseases from X-rays may be confident on common conditions but uncertain about rarer cases. By focusing human annotation on these ambiguous cases, active learning optimizes the use of expensive expert time while accelerating model improvement.




Figure 9.5: The Active Learning Closed Loop: The selection cycle trades proxy-scoring compute for queries chosen from an unlabeled pool for oracle annotation. Uncertainty, diversity, and coverage criteria seek more useful information per annotation than undirected sampling.

The economic implications can be substantial in specialist domains, where annotation cost and turnaround may dominate the compute used to score candidates. These query strategies drive each iteration of the active learning loop in figure 9.5, and a simple budget comparison shows how active learning can reduce labeling cost and training work under stated assumptions.

Napkin Math 9.3: The active learning ROI

Problem: A hospital team has 1 Million unlabeled scans, a \$5/label specialist-label cost, a \$500,000 budget, and 1 month. How does uncertainty sampling change labeling cost and per-epoch training work relative to the budget-matched random baseline?

Scenario A: Naive Labeling

- **Cost:** Labeling all 1 Million scans would cost \$5,000,000 (10× over budget).

- **Budget-matched baseline:** $\$500,000 \div \$5/\text{label} = 100,000$ random scans.
- **Naive labeling outcome:** Random selection does not guarantee rare-pathology coverage.

Scenario B: Active Learning

- **Strategy:** Uncertainty sampling selects 50,000 scans for specialist labeling.
- **Cost:** $50,000 \times \$5/\text{label} = \$250,000$ (50 percent under budget).
- **Training work:** $100,000 \div 50,000 = 2\times$ fewer examples per epoch than the budget-matched random baseline.
- **Active learning outcome:** Matching 100,000 random scans requires empirical validation.

Systems insight: Against the budget-matched random baseline, selection saves $\$500,000 - \$250,000 = \$250,000$; accuracy still requires empirical validation.

Targeted sample selection can reduce the annotation volume required to reach a target metric. Compare the two illustrative accuracy trajectories in figure 9.6, noting the horizontal gap between active uncertainty sampling and baseline random selection as sample counts scale.




Figure 9.6: The Active Learning Multiplier: Model accuracy is plotted against labeled samples on a logarithmic axis. In this illustrative scenario, active learning (green solid) reaches the same accuracy with fewer labels than random sampling (gray dashed); the arrow marks an approximately 4× label-efficiency gap near 90 percent accuracy.

The figure tracks a different axis from the calculation in napkin math 9.3. It illustrates label efficiency, while the callout calculates labeling dollars under separate assumptions. Active learning also incurs scoring, acquisition, and repeated-retraining costs, so end-to-end savings can be smaller than the label-count gap.

Active learning can do more than save labeling cost: a suitable query rule can concentrate review on uncertain, diverse, or failure-relevant examples.



Example 9.2: Hard negative mining in a smart doorbell

Scenario: A smart doorbell vision detector receives continuous video feeds dominated by empty backgrounds (easy negatives) and normal pedestrian poses.

Diagnosis: Random frame sampling can miss rare false positives such as statues, laundry piles, and human-shaped shadows. Active learning routes uncertain predictions (“Person: 51%”) and sampled production errors to human reviewers for targeted labeling.

Systems lesson: Hard-negative mining can improve labeling efficiency by spending review effort on suspected false positives and other difficult negatives. The query stream still needs diversity and coverage checks so confident but systematically wrong cases are not excluded.

15 | Illustrative clinical labeling economics: The review-rate and hourly-rate ranges are scenario assumptions rather than reported measurements. Under those inputs, labeling 500 scans requires 7–10 hours and costs \$1,000–3,000; labeling all 50,000 costs \$94,000–300,000. The arithmetic illustrates why semi-supervised learning can be attractive when expert labels are expensive, but actual review time, adjudication, prevalence, and validation requirements must be measured.

16 | Pseudo-labeling: Uses a trained model's own confident predictions as ground-truth labels for unlabeled data; the technique's effectiveness depends on a virtuous cycle: accurate predictions on easy unlabeled examples expand the training set, improving the model, which enables accurate predictions on harder examples. The failure mode is equally self-reinforcing: incorrect pseudo-labels reinforce errors through *confirmation bias*, making the confidence threshold a critical systems parameter. Setting it too low compounds label noise across training iterations; setting it too high wastes unlabeled data that could contribute learning signal.

17 | Consistency regularization: Rooted in the smoothness assumption: if two inputs x_1 and x_2 are close in input space, their labels should also be close; the training objective minimizes divergence between a model's predictions on an input and its augmented version. This is conceptually distinct from data augmentation (which creates more training examples) because it explicitly enforces prediction consistency as a loss term, even for unlabeled data where the "correct" label is unknown. The systems consequence is additional computation for generating and comparing predictions under multiple perturbations, a trade-off that can favor compute-rich, label-poor settings.

9.3.3 Semi-supervised learning: Using unlabeled data

Consider an illustrative medical-imaging scenario with 50,000 chest X-rays, of which 500 have been reviewed and labeled¹⁵ by radiologists, a labeling rate of 1 percent. Whether that seed set is adequate depends on the task and distribution. Semi-supervised learning can use the remaining 49,500 images when the unlabeled pool matches the target distribution and the method is validated on labeled data.

Active learning optimizes which samples to label but still requires human annotation for every selected example. Semi-supervised learning takes a more aggressive approach by extracting learning signal from unlabeled data. It uses a labeled subset to guide learning on a larger unlabeled pool and can approach fully supervised accuracy with substantially fewer labels, although the result depends on the task, distributions, and method.

The core insight behind semi-supervised learning is that unlabeled data, while it cannot directly teach the mapping from inputs to outputs, contains structural information about the input distribution $p(x)$ that can constrain the hypothesis space. Under smoothness and cluster assumptions, a decision boundary that cuts through a dense region assigns different labels to nearby inputs and is less plausible than one in a low-density region. Semi-supervised methods exploit these assumptions, which must be checked against the target task.

Three main techniques implement this insight. Pseudo-labeling¹⁶ takes the most direct approach: train on labeled data, use the model to generate "pseudo-labels" for high-confidence unlabeled predictions, then retrain on both. The confidence threshold is critical: setting it too low introduces label noise that degrades learning, while setting it too high wastes potentially useful data.

Consistency regularization¹⁷ takes a different angle by enforcing that the model produces similar predictions for augmented versions of the same input. A robust classifier should be invariant to realistic perturbations like cropping, rotation, or color shifts. Methods like FixMatch¹⁸ combine both approaches, assigning pseudo-labels only to samples where the unaugmented prediction is confident but training the model to predict these labels on strongly augmented versions of the same images.

Label propagation offers a third paradigm through graph-based reasoning: construct a similarity graph over all samples and propagate labels from labeled nodes to their neighbors. This approach works particularly well when the feature space exhibits clear cluster structure.

Example 9.3: FixMatch on CIFAR-10

Scenario: A semi-supervised image classification pipeline trains on CIFAR-10 with only 250 labeled samples (25 per class) alongside unlabeled data (Sohn et al. 2020).

Diagnosis: The fully labeled CIFAR-10 training set contains 50K labels. FixMatch generates high-confidence pseudo-labels (>0.95 threshold) on weakly augmented unlabeled images and enforces consistency on strongly augmented variants, reaching 94.9 percent accuracy in the cited experiment.

Systems lesson: Semi-supervised learning trades additional training compute for fewer manual labels. Under the chapter's labeling-cost assumptions, the FixMatch scenario reduces dataset preparation expense by 8.1×; the end-to-end saving depends on the unlabeled-data and compute costs.

The systems trade-off in semi-supervised learning exchanges labeling effort for additional computation over labeled and unlabeled samples. It can approach fully supervised accuracy with fewer labels, but neither the label reduction nor the compute increase is universal. A CIFAR-10 comparison makes one such trade-off concrete (table 9.6).

The efficiency gains are substantial, but semi-supervised learning is not universally applicable. The technique assumes that unlabeled data comes from the same distribution as labeled data, and it

struggles when unlabeled data contains out-of-distribution samples (the model confidently mislabels them), when class imbalance is severe (pseudo-labels amplify majority class bias), or when the labeled set does not cover all classes (preventing label propagation for unseen classes). Always validate on a held-out set with true labels to catch distribution mismatch.

Table 9.6: FixMatch Label Efficiency on CIFAR-10: With 250 labels (0.5 percent of the dataset), FixMatch reaches 94.9 percent accuracy while using 200× fewer labels than the full training set contains.

Label Budget	Method	Accuracy	Label Efficiency
50,000 (100%)	Full training set	—	Label-count reference
4,000 (8%)	FixMatch	95.7%	12.5× more efficient
250 (0.5%)	FixMatch	94.9%	200× more efficient
40 (0.08%)	FixMatch	88.6%	1250× more efficient

Despite these limitations, semi-supervised learning can reduce label requirements while maintaining accuracy on suitable tasks. Across the techniques, coreset selection and deduplication prune low-value samples before training; curriculum learning changes the presentation order; active learning chooses which samples receive human annotation; and semi-supervised learning extracts additional signal from an unlabeled pool. Each technique can reduce dependence on task-specific labels, but none eliminates the need for labeled evaluation. The progression raises a related possibility: task-specific labels may not be necessary for pretraining. The structure of data itself—that cat images resemble other cat images and coherent sentences follow grammatical patterns—may provide a supervision signal, while downstream quality still requires task-aligned validation.

9.4 Self-Supervised Learning

GPT was trained to predict the next token in a sequence. BERT was trained to fill in masked tokens. Neither pretraining objective required task-specific human labels. **self-supervised learning**¹⁹ generalizes this insight: by designing *pretext tasks* that derive supervision from the data’s structure, models can learn reusable representations from unlabeled data at scale. Where active and semi-supervised learning reduce the demand for labels, SSL changes the accounting by making unlabeled structure the pretraining signal. In the three-stage map, this makes SSL less a fourth selection stage than an extension of the available pool: the question shifts from which labeled samples to train on to how unlabeled samples can support pretraining. It responds to one aspect of the data wall in section 9.1.4 by broadening what counts as a training signal, while leaving corpus quality, licensing, coverage, and labeled evaluation requirements intact.

Labels are one form of supervision; a pretext task can also derive a learning signal from the structure of the data, as table 9.7 summarizes.

Table 9.7: Self-Supervised Pretext Tasks by Modality: Each task extracts supervision from data structure rather than human labels, enabling pretraining on large-scale unlabeled corpora.

Modality	Self-Supervised Task	Supervision Signal
Text	Masked language modeling	Predict [MASK] from context
Text	Next-token prediction	Predict next token in sequence
Images	Contrastive learning	Same image (augmented) vs. different images
Images	Masked autoencoding	Reconstruct masked patches
Multi-modal	CLIP-style alignment	Match image-text pairs

Pretext tasks generate supervision signals automatically. Masked-token objectives can encourage models to learn grammatical and semantic structure; next-token prediction teaches continuation structure; and contrastive image objectives encourage visual features that are invariant to selected transformations. These approaches build on the convolutional neural network and transformer families in chapter 6, but the systems accounting matters: self-supervision removes task-specific

18 | FixMatch: It generates a pseudo-label using a weakly augmented image and then trains the model to predict that same label for a strongly augmented version of the image. This consistency training is gated by a confidence threshold; a pseudo-label is only used if the model’s prediction on the weak augmentation is highly confident (for example, >0.95) (Sohn et al. 2020). This embodies a direct systems trade-off between additional training compute and fewer manual labels; the magnitude depends on the baseline and cost model.

19 | Self-supervised learning: The pretext task, such as next-token or masked-token prediction, provides a supervisory signal derived from the data itself. This reduces dependence on task-specific human labels but can require substantial pretraining compute. Amortization depends on how broadly the resulting model is reused.

annotation from the pretraining gate rather than removing data cost. Pretraining can begin on unlabeled corpora without waiting for example-level labels, but collection, filtering, licensing, storage, and processing remain. Separating pretraining from downstream labeling restructures the economics of machine learning when the representation transfers across tasks; task-specific adaptation and evaluation still carry their own data requirements.

9.4.1 The economics of amortization

Understanding *why* self-supervised learning became central to many foundation-model workflows requires examining its economic structure. A foundation model is a broadly pretrained reusable base model that can be adapted to many downstream tasks through fine-tuning. That shift translates into concrete cost savings through *cost amortization*, where expensive pretraining is performed once and reused across many applications (table 9.8).

Table 9.8: Illustrative Foundation-Model Amortization: The label and compute ranges are scenario assumptions. They illustrate why reuse can lower marginal task cost when representations transfer, not a universal fine-tuning ratio.

Approach	Labels per Task	Compute per Task	Data Acquisition
Train from scratch	100K–1M labeled	100% full training	Task-specific collection
Fine-tune foundation model	100–1K labeled	1–5% of full training	Reuse pretraining corpus

To illustrate this economic transformation, consider a company building 10 specialized classifiers. Under the scenario assumptions, each task trained from scratch needs 100,000 labels at \$1 per label, for \$1M across all tasks. The modeled compute burden is 10,000 GPU-hours.

The fine-tuning scenario pays 10,000 GPU-hours once, then assigns 1,000 labels and 50 GPU-hours of compute to each task. Across all 10 tasks, the modeled fine-tuning labels cost \$10K.

Under these inputs, labeling cost drops by 100× and per-task marginal compute by 20×. These are consequences of the chosen scenario values, while total compute becomes favorable only after enough downstream tasks amortize pretraining.

The cost structure explains why fine-tuning can be economical when a pretrained model transfers across many downstream applications. The result depends on pretraining cost, reuse count, and per-task adaptation requirements.

Investing in foundation model pretraining fundamentally alters downstream task economics. Compare the side-by-side cost bar charts in figure 9.7, contrasting repetitive scratch training expenses on the left against amortized pretraining with low-cost fine-tuning on the right.




Figure 9.7: Cost Amortization in Foundation Models: Training from scratch (left) requires 1,000 GPU-hours per task (10,000 GPU-hours total for 10 tasks). The foundation model approach (right) pays 10,000 GPU-hours upfront for pretraining but reduces each subsequent task to just 50 GPU-hours. At 10 tasks the totals are comparable (10,000 GPU-hours vs. 10,500 GPU-hours), but the per-task marginal cost drops by 20×, and the crossover favoring the foundation model occurs around 10.5 tasks. The plotted 10-task snapshot therefore remains just below that break-even point.

9.4.2 Foundation model paradigm

The amortization economics can favor self-supervised learning, though different SSL methods occupy different points on the cost-efficiency frontier. The original SimCLR experiments (T. Chen et al. 2020) benefited from batches as large as 4,096, while MoCo (He et al. 2020) uses a queue-based dictionary and momentum encoder to decouple the number of negatives from the mini-batch. Masked modeling methods have different memory and step-count trade-offs. Generative pretraining follows empirical power-law scaling with data, parameters, and compute, making it attractive when its cost is amortized across downstream tasks. These methods rely on architecture families introduced in chapter 6; what matters for data selection is that self-supervised pretraining can increase the value of limited labeled data when the learned representation transfers.

When its pretraining cost is reused across tasks, SSL supports the **foundation model paradigm**²⁰ (Bommasani et al. 2021). The data selection principles discussed throughout this chapter (coreset selection, curriculum learning, active learning) remain relevant within this paradigm. Pretraining corpus curation applies the same deduplication and quality filtering techniques at web scale, and fine-tuning data selection determines which labeled examples maximize downstream task performance.

Self-supervised learning addresses the label bottleneck by learning from data structure rather than human annotation, yet it cannot solve data scarcity itself. Rare classes may have too few examples, edge cases may never appear in the wild, and privacy constraints may prevent collecting real samples. The third stage of the data selection pipeline addresses this gap by creating new data on demand rather than selecting or curating existing data.

9.5 Synthetic Data Generation

A robot fleet may need rare collision examples, a medical model may need cases a hospital cannot share, and a wake-word model may need thousands of noisy kitchens before the product exists. Static pruning removed redundancy before training. Dynamic selection focused compute on the most informative samples during training, and self-supervised pretraining drove its label requirement to zero, reframing what counts as training data rather than adding a stage. The third and final stage of the data selection pipeline, synthetic data generation, takes the opposite approach: rather than subtracting or selecting from existing data, it creates new high-value samples when real data is scarce, expensive, or lacks diversity. The strategy shifts from *curation* to *creation*.

9.5.1 Data augmentation: Transformation-based synthesis


Data augmentation is the lowest-cost form of synthetic generation: it expands a dataset by applying transformations to existing samples. Because many transformations preserve label semantics while creating novel inputs, augmentation effectively multiplies the diversity of a training set without requiring additional data collection.

For image data, augmentation should match the invariance the deployed model must learn (Shorten and Khoshgoftaar 2019). Geometric transformations such as rotation, flipping, cropping, and scaling introduce spatial variation that makes models robust to viewpoint changes. Photometric transformations adjust brightness, contrast, saturation, and hue to simulate different lighting conditions and camera characteristics. More advanced techniques like Cutout²¹ (which applies random rectangular masks), MixUp (H. Zhang et al. 2018) (which blends two images and their labels), and CutMix²² (which pastes patches between images) push augmentation further by creating entirely synthetic training examples that regularize learning.

Text augmentation presents different challenges because language is discrete rather than continuous. Back-translation²³ offers one solution for machine translation by translating target-language monolingual text into the source language and pairing the synthetic source with the original target. Simpler approaches include synonym replacement and random insertion or deletion, although these transformations can change meaning and must be validated for the task.

Rather than hand-designing these augmentation policies, AutoAugment²⁴ uses reinforcement learning to discover augmentation strategies for specific datasets, while RandAugment²⁵ simplifies this by sampling from a fixed set of transformations. That compute-quality trade-off depends on both the target architecture and the invariances required at deployment.

²⁰ | **Foundation model:** The name emphasizes that these models serve as a shared base for many downstream tasks, but this creates a single point of failure. Defects in the foundation model's pretraining data (biases, factual errors, memorized private content) propagate to every application built upon it. From a systems perspective, this homogenization risk means that data selection quality during pretraining has an out-sized blast radius: a curation error that would affect one task in the train-from-scratch paradigm now affects thousands of downstream deployments.



One pretraining corpus defect propagates into many downstream tasks.

²¹ | **Cutout:** Randomly masks square regions of input images during training, forcing the model to recognize objects from partial information rather than relying on any single discriminative region (DeVries and Taylor 2017). Unlike dropout (which zeroes neurons in feature space), Cutout operates in input space. The original Cutout experiments produced 0.3–2.0 percentage-point gains across CIFAR-10/100 and SVHN with negligible compute overhead, making it a high-ICR augmentation technique when occlusion-style invariance matches the task: more information per sample at near-zero additional cost to the pipeline.

22 | CutMix: Replaces Cutout’s zeroed-out region with a patch from a different training image, mixing labels proportionally to patch area (30 percent of image A replaced by image B yields a 70/30 label split) (Yun et al. 2019). This addresses Cutout’s weakness: zeroed regions waste pixel information that could carry learning signal. The CutMix paper reports ImageNet top-1 improvements of 2.28 percentage points for ResNet-50 and 1.70 percentage points for ResNet-101, while also improving localization behavior. The method provides stronger regularization than either Cutout or MixUp alone with little additional data-pipeline cost.

23 | Back-translation: In the method described by Senrich et al. (2016), target-language monolingual text is translated into the source language, and the synthetic source is paired with the unchanged target for machine-translation training. The systems trade-off is latency because each synthetic source requires translation-model inference. Pipelines can precompute these examples offline rather than generating them in the data loader.

24 | AutoAugment: Treats augmentation policy design as a reinforcement learning search problem: a controller selects operations (rotate, translate, shear, equalize), their magnitudes, and application probabilities to maximize validation accuracy (Cubuk et al. 2019). The original ImageNet search used approximately 15,000 GPU-hours, motivating later methods that reduce the search space.

25 | RandAugment: Collapses AutoAugment’s policy search to two hyperparameters: transformation count N and shared magnitude M (Cubuk et al. 2020). That small search space matches or exceeds AutoAugment on the reported benchmarks at far lower search cost, making it practical when 15,000 GPU-hours of policy search cannot be justified.



Lighthouse 9.2: MobileNetV2 and aggressive augmentation

Our MobileNetV2 lighthouse model from section 6.1.1 exemplifies how data augmentation can complement a resource-constrained architecture. Its depthwise separable convolutions reduce computation and parameter count relative to comparable standard convolutions, while augmentation can improve generalization when training data is limited.

The appropriate augmentation strength depends on the dataset, model, and invariances required at deployment. Techniques such as cropping and RandAugment can increase training diversity without increasing inference-time model capacity, but their magnitude and resulting accuracy must be measured for the target workload.

9.5.2 Generative synthesis: Creating new samples

Augmentation transforms existing samples; synthetic data generation goes further by creating new examples using generative models or simulators. It can be useful in three common scenarios:

- **Privacy-sensitive data:** Synthetic examples can reduce direct exposure to real records, but still require privacy auditing for memorization or leakage.
- **Rare edge cases:** Synthetic examples can cover failure modes, such as autonomous driving scenarios, that must be tested but seldom occur naturally.
- **Expensive collection:** Synthetic examples can supplement domains such as robotics or scientific experiments where each real sample requires physical resources.

The distribution sketch in figure 9.8 shows the central risk of that strategy: synthetic examples can be plentiful and still misaligned with deployment data.




Figure 9.8: The Domain Gap Problem: Synthetic data (blue) and real data (orange) occupy shifted feature distributions. The dashed red synthetic boundary and dotted green real-data boundary illustrate how a boundary learned for one distribution can be misaligned with the other.

The choice among generative approaches is a cost-fidelity decision. Generative adversarial networks (GANs) train a generator against a discriminator in an adversarial training setup, producing realistic images through competition; StyleGAN, for instance, generates photorealistic faces that have augmented facial recognition datasets. Diffusion models use iterative denoising to produce high-quality images; systems like Stable Diffusion²⁶ support targeted training example generation from natural language descriptions through text-to-image synthesis. Finally, simulation engines such as CARLA for autonomous driving or Unity and Unreal for robotics offer physics-based rendering that can generate large volumes of labeled scenarios with known simulator state, making them particularly valuable for safety-critical applications where edge case coverage is essential.

9.5.3 Bridging the domain gap

Relying on synthetic data generation introduces severe distribution alignment risks.²⁷ Examine the feature space mapping in figure 9.8, noting how decision boundaries trained on synthetic data shift away from real production deployment clusters.

Two complementary strategies address this distribution mismatch. Domain randomization²⁸

takes an aggressive approach: rather than trying to match the real world precisely, it trains on wildly varied synthetic data by randomizing lighting, textures, backgrounds, and camera parameters during generation.

The aim is to make deployment conditions fall within the variation seen during training, but coverage is never guaranteed by randomization alone. The method can be effective in robotics and autonomous-driving settings when the simulator spans the relevant physical and visual factors and the result is validated on real data.

Domain adaptation takes the opposite approach by explicitly aligning synthetic and real distributions. Feature alignment methods train on synthetic data while simultaneously minimizing the distance between synthetic and real feature distributions, often using adversarial training to learn domain-invariant representations. Fine-tuning offers a simpler path: pretrain on abundant synthetic data to learn general features, then fine-tune on a small real dataset to adapt to deployment conditions. Self-training combines these ideas by using a synthetic-trained model to pseudo-label real unlabeled data, then retraining on the combined labeled set.

In practice, mixing synthetic and real data provides a real-data anchor while retaining some synthetic coverage, but the best ratio is workload-dependent. Table 9.9 summarizes the trade-off across different mixing ratios.

Table 9.9: Synthetic-to-Real Data Mixing Ratios: Pure synthetic data can suffer from distribution shift; pure real data is expensive. The optimal ratio varies by domain, simulator fidelity, model family, and validation protocol, so the table should be read as a scenario scaffold rather than a universal recipe.

Synthetic Fraction	Representative Outcome
100% synthetic	No real-data anchor; validate the domain gap
80% synthetic + 20% real	Lower real-data demand; validate on deployment data
50% synthetic + 50% real	More real-data anchoring; higher collection cost
100% real	No synthetic-to-real gap; highest collection cost

 Example 9.4: KWS data selection

Scenario: An embedded keyword spotting (KWS) system targets deployment on a 256 KB SRAM microcontroller with a 10,000+ audio sample requirement.

Diagnosis: Manually recording 10,000 real utterances across diverse speakers and acoustic environments costs \$20K–50K. Staging a 5-tier pipeline (seed recordings → augmentation → noise injection → hard-negative mining → TTS simulation) builds dataset scale from 500 seed samples.

Systems lesson: Layered data strategies can lower acquisition costs for resource-constrained edge ML deployments. In this scenario, augmentation and noise injection reduce physical collection cost to 5 percent of the full-recording estimate; target accuracy and acoustic coverage still require validation.

The keyword-spotting scenario shows how augmentation, noise injection, and simulation can expand a small seed set, but the same creation strategy becomes riskier when samples come from learned generators. In recursive training, *model collapse*²⁹ can amplify errors and reduce diversity over generations. Synthetic data therefore requires provenance, mixture controls, and validation against real deployment data rather than a blanket assumption that more generated samples help.

9.5.4 Knowledge distillation: Compressing information

While data augmentation (section 9.5.1) and generative synthesis (section 9.5.2) create new input samples, another form of synthesis creates teacher-generated targets. **Knowledge distillation**³⁰ (Hinton et al. 2015; Gou et al. 2021) trains a student model using a teacher model’s outputs alongside or instead of hard labels. Soft targets can expose relative class probabilities that a one-hot label omits, although their value depends on teacher quality, temperature, and task alignment.

26 | **Stable Diffusion:** Performs iterative denoising in a compressed latent space, enabling targeted text-to-image examples (Rombach et al. 2022). Its iterative generation is substantially more compute-intensive than geometric transforms, so it fits cases where novelty matters more than per-sample throughput.

27 | **Domain gap:** The statistical divergence between synthetic and real data distributions, measurable with metrics such as maximum mean discrepancy (MMD) or Frechet Inception Distance (FID). For ML systems, the gap becomes training-serving skew: a simulator-trained model can validate well on synthetic data while failing silently on real deployment data.

28 | **Domain randomization:** Makes synthetic data deliberately varied by randomizing textures, colors, lighting, and physical properties. The goal is not photorealism; it is enough variation to improve coverage of deployment conditions, shifting the cost bottleneck from rendering fidelity to coverage.

29 | **Model collapse:** Formally analyzed by Shumailov et al. (2024), this phenomenon occurs because generative models systematically under-represent tail distributions – rare but important patterns that appear infrequently in training data. When generation $n + 1$ trains on output from generation n , each successive generation further compresses the tails, producing increasingly homogeneous data. The degradation can be rapid in recursive training settings, which is why the Fallacies section of this chapter warns against pure synthetic training.

30 | Knowledge distillation: Soft teacher probabilities can reveal relationships among classes that one-hot labels omit. The temperature parameter controls how soft those probabilities become; it must be tuned because excessive smoothing can erase useful distinctions. Distillation adds teacher inference and does not guarantee a higher ICR unless the transferred signal improves the student enough to justify that cost.

A teacher prediction such as [0.7, 0.2, 0.1] exposes relative preferences that a one-hot label [1, 0, 0] omits. A student may benefit from that signal when the teacher is accurate and aligned with the target distribution. At scale, teacher-generated targets add an inference pass and require quality controls; any student cost or quality advantage must be measured against that added generation expense.

Together, augmentation, generative synthesis, and distillation complete the third stage of the data selection pipeline. Where static pruning removes redundancy and dynamic selection focuses compute on high-value samples, synthetic generation fills gaps by creating samples that never existed. Selecting the right combination of techniques from these three pipeline stages requires a structured decision framework.

9.6 Decision Framework

When labeling budget, redundancy, rare classes, privacy, and convergence speed all matter at once, the practitioner first has to identify which constraint is binding. Each stage can raise the information-compute ratio when its measured quality gain justifies its overhead: pruning removes low-value samples before training, dynamic selection focuses compute on high-value samples during training, and synthesis creates new high-value samples on demand (table 9.10).

Table 9.10: Three-Stage Data Selection Pipeline: The ranges illustrate possible outcomes rather than portable expectations. Actual gains depend on the dataset, model, selection method, and target metric.

Stage	When Applied	Techniques	Illustrative Outcome
Static pruning	Before training	Coreset Selection, Deduplication, Quality Filtering	30–50% dataset reduction
Dynamic selection	During training	Curriculum Learning, Active Learning, Semi-Supervised	10–30% faster curricula; 2–100× fewer labels
Synthetic generation	On-demand	Augmentation, Generative Models, Distillation	2–10× effective data expansion

Once the dominant constraint is named, table 9.11 compares which technique changes that constraint and why.

Table 9.11: Technique Selection Guide by Primary Constraint: Candidate data-selection techniques mapped to the dominant resource constraint in the ML pipeline. The table begins the diagnosis; prerequisites and measured ROI determine the final choice.

Constraint	Candidate Technique	Why
Limited labeling budget	Active Learning	Maximizes label ROI by selecting informative samples
High redundancy in data	Deduplication + Coreset	Removes waste before training begins
Rare classes or edge cases	Synthetic Generation	Creates samples that do not exist in raw data
Slow convergence	Curriculum Learning	Improves gradient quality in early training
Privacy requirements	Privacy-audited synthesis	Reduce direct use of real records; verify leakage
Large model, small dataset	Knowledge Distillation	Use teacher model’s knowledge as “data”

Table 9.11 maps individual constraints to techniques, but real projects face multiple constraints simultaneously. The decision tree in figure 9.9 structures the selection process hierarchically: start by identifying the primary bottleneck, then follow the branches to narrow the field.

Decision process

Each path requires a structured assessment because the same dataset symptom can point to different bottlenecks. Technique selection begins by naming the binding constraint, then checks whether the data, labels, and infrastructure needed by the chosen method exist.




Figure 9.9: Data Selection Decision Tree: The tree begins with the primary bottleneck—labeling cost, compute cost, or data scarcity—and narrows each branch using available pools, oracles, redundancy, teachers, or simulators. Its leaves identify candidate techniques; projects with multiple binding constraints may require more than one branch.

Step 1: Assess the bottleneck

Identify which resource constraint most severely limits the training pipeline:

- **Labeling cost:** Label-efficiency techniques such as active learning, semi-supervised learning, and self-supervised learning maximize the value extracted from each human annotation.
- **Compute cost:** Coreset selection and deduplication can reduce processed samples, while a validated curriculum may reduce the steps needed to reach a target metric.
- **Data scarcity:** Data creation through augmentation, synthesis, and distillation expands the effective training set beyond what raw collection provides.

This diagnostic step keeps technique selection tied to the binding resource constraint rather than to the most familiar algorithm.

Step 2: Check prerequisites

With the bottleneck identified, verify that the corresponding techniques are feasible given the available infrastructure and data. Each approach carries specific requirements that must be met before implementation can begin (table 9.12).

Table 9.12: Technique Prerequisites: Each technique carries specific infrastructure and data requirements that must be verified before implementation. A technique with excellent theoretical gains but unmet prerequisites will fail in practice, making this checklist the first step in technique selection.

Technique	Prerequisites
Active Learning	Access to oracle, unlabeled pool, retraining infrastructure
Coreset selection	Proxy model or embedding extractor, full dataset accessible
Curriculum Learning	Difficulty scoring method, pacing schedule
Semi-Supervised	Some labeled data, unlabeled data from same distribution
Self-Supervised	Large unlabeled corpus, pretraining compute budget
Augmentation	Domain knowledge of invariances, augmentation library
Synthetic Generation	Generative model or simulator, domain gap mitigation

Step 3: Estimate ROI

Before committing engineering resources, estimate each candidate technique’s return on investment:

ROI =
$$\frac{(\text{Baseline Cost}) - (\text{Technique Cost} + \text{Implementation Cost})}{\text{Technique Cost} + \text{Implementation Cost}}$$

A technique with high theoretical gains but high implementation cost may deliver lower ROI than a simpler approach. Exact deduplication is often an inexpensive first candidate because hashing is simple and repeated bytes are easy to identify, although downstream benefit still depends on corpus redundancy and intentional sample weighting. Active learning requires oracle access, retraining infrastructure, and selection algorithm development, so its ROI depends heavily on how many labeling cycles amortize that investment.

Step 4: Combine techniques

The techniques in this chapter are not mutually exclusive, but combining them does not guarantee additive gains. One possible workflow deduplicates the raw corpus, applies a coverage-aware coreset method, orders selected samples with a validated curriculum, and augments them at runtime. A pretrained representation can supply a different starting point from random initialization. Each added stage must earn its place because overlapping filters can remove the same examples while their overheads continue to accumulate.

Stages can compound efficiency gains when their effects do not overlap and their overhead remains small; combined savings must be measured rather than assumed.

This decision framework answers the *what* of data selection: which samples to prune, when to select dynamically, and how to synthesize new data. Understanding these algorithmic choices is essential, but algorithms alone do not translate into faster training. A perfectly designed coreset algorithm that takes 10 hours to select samples for a two-hour training run yields no practical benefit. Similarly, a curriculum learning strategy that requires scanning the entire dataset to determine difficulty rankings may idle GPUs while CPUs compute scores. The *how* of implementation matters as much as the *what* of algorithm choice.

The gap between algorithmic elegance and practical value raises several systems challenges: preventing selection overhead from negating theoretical gains, handling nonsequential I/O patterns that confuse prefetching logic, and coordinating selection decisions across distributed workers without introducing synchronization bottlenecks. The engineering patterns that follow bridge the gap between data selection theory and production reality.

9.7 Selection Engineering

A naive active learning loop that scans the entire dataset every epoch to select the “best” samples can turn a compute-bound training job into an I/O-bound bottleneck. **Selection engineering** begins where the decision framework ends: after identifying which algorithms to apply, engineering must ensure that they deliver their promised speedups on real hardware and real data pipelines. The architectural patterns that prevent this failure implement data selection in production.

9.7.1 The selection bottleneck

Dynamic data selection introduces a new bottleneck: selection latency. A conventional loader follows a predetermined sampling plan, whereas active learning or adaptive curricula may evaluate a selection function $f(x)$ over a large candidate pool before forming later batches. Concretely, scoring a 1M dataset with a large model can take 2.8 hours, potentially negating the savings from a 10 percent coreset if not performed with a smaller proxy model. The systems trade-off requires that selection cost remain below the training work the subset saves; otherwise, sample counts fall while end-to-end cost rises.

At a fixed target metric, selection plus subset training must cost less than full-data training. Let $T_{\text{selection}}$ be wall-clock time spent scoring the pool, $T_{\text{train}}(D)$ be wall-clock training time for a dataset of size D , and D_{subset} and D_{total} be the retained and full sample counts. Equation 9.2 expresses this break-even condition, the **selection inequality**.

$$T_{\text{selection}} + T_{\text{train}}(D_{\text{subset}}) < T_{\text{train}}(D_{\text{total}}) \quad (9.2)$$

If $f(x)$ repeatedly scores the pool with a large model, recurring rescoring and coordination can consume the training savings and produce negative ROI. A concrete scenario illustrates this trade-off.

Equivalently, we can isolate the allowable selection overhead: $T_{\text{selection}} < T_{\text{train}}(D_{\text{total}}) - T_{\text{train}}(D_{\text{subset}})$. The allowable overhead therefore depends on how much training work the subset saves; no fixed percentage applies across workloads.



Example 9.5: Selection inequality in practice

Scenario: A computer vision team selects a 100K coreset (10 percent) from 1M images to reduce training time on an accelerator cluster.

Diagnosis: Full-model scoring (Option A: target ResNet-50) adds 2.8 h of selection overhead, reducing net savings to 247.2 h. Proxy-model scoring (Option B: ResNet-18) drops selection overhead to 0.6 h, saving 249.4 h total wall-clock hours.

Systems lesson: Coreset selection reduces training latency only when selection overhead is smaller than full-dataset training savings. Proxy scoring preserves net FLOP savings by keeping sample scoring overhead well below full-model forward pass costs.

Data selection improves end-to-end efficiency only if its overhead is smaller than the training time saved on the reduced dataset. Compare the three illustrative stacked time bars in figure 9.10, observing how an expensive selection function can consume all subset-training savings.




Figure 9.10: The Selection Inequality: Data selection only improves end-to-end efficiency if the overhead of selection plus training on the subset is less than training on the full dataset. A lightweight selection function (proxy model, cached embeddings) keeps selection overhead low; an expensive selection function (full model forward pass) can negate the savings.

The consequence shown in figure 9.10 is unambiguous. Selection overhead can negate the benefits of training on a smaller subset. The cost compounds when selection runs every epoch instead of once; if per-epoch scoring approaches the training work saved by the subset, the net advantage disappears.

9.7.2 Hardware empathy: The random-access penalty

The selection inequality addresses compute overhead, but some selection strategies also change I/O patterns. Index-driven sampling can turn large sequential shard reads into smaller scattered reads, reducing the benefit of readahead and batching. The impact depends on record size, queue depth, cache state, storage tier, and data-loader design. Table 9.13 illustrates one 4 KB, single-device comparison rather than a training-pipeline benchmark. Selection cost therefore includes the access pattern it induces. Equal-sized subsets can produce different wall-clock savings when one preserves sequential reads and the other fragments them.

Systems can mitigate this penalty with proxy models, cached embeddings, and indexed retrieval. A smaller proxy can reduce scoring cost only if its ranking transfers to the target model. Lower-precision inference may reduce that cost further when the runtime, model, and required ranking fidelity support it. Embedding indices such as FAISS³¹ replace some full scans with exact or approximate retrieval whose cost and recall depend on the index configuration. These approaches decouple selection from target-model training so each can be measured independently.

³¹ | **FAISS (Facebook AI Similarity Search):** Provides GPU-accelerated similarity search using exact, approximate, and compressed-domain index designs (Johnson et al. 2019); Johnson, Douze, and Jegou report billion-vector graph construction on multiple GPUs, showing why vector-index infrastructure matters for web-scale selection. For data selection pipelines, FAISS-style indexing supports k -nearest-neighbor retrieval for coreset selection, embedding-based deduplication, and stratified clustering for balanced sampling. Without this infrastructure, embedding-based selection would often fall back to expensive full-corpus scans.

Table 9.13: Illustrative Small-Read Penalty: Registry device anchors compare sequential bandwidth with random 4 KB reads under a simplified IOPS conversion. Real throughput depends on request size, queue depth, concurrency, caching, and service configuration; the cloud row is only a qualitative scenario.

Storage Tier	Sequential Throughput	Random I/O (IOPS)	Random Throughput (approx)	Random Penalty
HDD (7.2k)	~150 MB/s	~100 IOPS	~0.4 MB/s	375×
SATA SSD	~550 MB/s	~10K IOPS	~40 MB/s	13.8×
NVMe SSD	~3,500 MB/s	~500K IOPS	~2,000 MB/s	1.75×
Cloud (S3)	Configuration-dependent	Not IOPS-comparable	Configuration-dependent	Potentially extreme

Data loaders also require architectural adaptation. Sharded dataset formats package many samples into sequentially readable chunks; WebDataset and FFCV are concrete examples. Shuffle buffers are a data-machine co-design: the loader reads large sequential shards into memory and samples randomly within the buffer, preserving sequential I/O throughput while achieving the statistical benefits of random sampling. In multi-accelerator training, each worker maintains its own shuffle buffer over a non-overlapping shard, so the randomization is local rather than global; rare classes or boundary cases that appear in only a few shards receive uneven coverage across workers unless the coreset is first stratified and shards are balanced before distribution.



Checkpoint 9.2: The selection inequality

Data selection is not free. It introduces a new term to the iron law and a new I/O cost.

Equation checks:

- ☐ **Selection cost:** can you explain why $T_{\text{selection}} + T_{\text{train}}(\text{subset})$ must be less than $T_{\text{train}}(\text{full})$ for the technique to reduce end-to-end training time?
- ☐ **Overhead management:** how do proxy models and cached embeddings keep $T_{\text{selection}}$ low?

Systems implications:

- ☐ **I/O patterns:** why can nonsequential access for dynamic selection reduce data-loader throughput relative to sequential reads?

9.7.3 Data echoing: Amortizing I/O costs

The optimizations discussed so far address I/O bandwidth, but modern data selection pipelines introduce another bottleneck: CPU computation. Synthetic data generation and heavy augmentation shift the constraint from disk speed to augmentation throughput. Heavy augmentations like 3D rotations and MixUp, or on-the-fly generative synthesis, can leave the GPU idle if the CPU cannot keep pace with sample production. When the data pipeline produces samples slower than the GPU can consume them, GPU utilization drops and training time extends, negating the efficiency gains from smarter data selection.


Data echoing³² (Choi et al. 2019) reuses data multiple times before fetching new samples, trading freshness for accelerator utilization when the input pipeline is slower than training. Echoing before randomized augmentation can produce different transformed inputs on each repetition, while echoing later in the pipeline may repeat identical tensors.

The optimal echo factor depends on the ratio R of upstream processing time to downstream training time:

$$R = \frac{T_{\text{data pipeline}}}{T_{\text{GPU training}}}$$

If $R > 1$ (data pipeline is the bottleneck), an echo factor $e < R$ partially recovers idle GPU cycles, while $e \geq R$ can fully use GPU capacity if echoed samples remain statistically useful. Increasing e beyond R no longer improves utilization and can reduce sample diversity. If $R < 1$ (GPU is the bottleneck), data echoing provides no benefit. A realistic scenario makes these trade-offs concrete.

32 | Data echoing: The key subtlety is *where* in the pipeline to insert the echo point. Echoing before augmentation (upstream echoing) can apply different random augmentations to each repetition, while echoing after augmentation feeds identical tensors to the accelerator. Useful echo factors and realized speedups depend on the workload, batch size, insertion point, and shuffling policy (Choi et al. 2019).



Pipeline utilization saturates near the ratio threshold; statistical value remains workload-dependent.

Napkin Math 9.4: Worked example: Data echoing ROI

Scenario: Training ResNet-50 on ImageNet with heavy augmentation (RandAugment + MixUp).

Measurements:

- Data pipeline throughput: 300 images/s (reading, decoding, augmenting on CPU)
- GPU training throughput: 800 images/s (forward + backward pass)
- Ratio $R = T_{\text{pipeline}}/T_{\text{GPU}} = (1/300 \text{ images/s}) / (1/800 \text{ images/s}) = 800 \text{ images/s} / 300 \text{ images/s} \approx 2.67$ (GPU waiting 62.5 percent of time)

Without echoing:

- Effective throughput: 300 images/s (limited by data pipeline)
- Training time for 90 epochs: $90 \times 1.28\text{M} / 300 \text{ images/s} = 384,350 \text{ seconds}$ (106.8 hours)
- GPU utilization: 37.5 percent

With echo factor: $e = 2$.

- Each batch is processed twice with different augmentations
- Effective throughput: 600 images/s (still below GPU capacity)
- Unique images per second: 300 images/s (unchanged)
- Training time: $90 \times 1.28\text{M} / 600 \text{ images/s} = 192,175 \text{ seconds}$ (53.4 hours) if echoed data is equally valuable

Trade-off: Repeated samples are not guaranteed to be as useful as fresh samples; the data echoing paper evaluates echo factor, insertion point, and shuffling because those choices determine whether reuse preserves predictive performance. In one network-fed ResNet-50/ImageNet configuration, Choi et al. (2019) reports a $3.25\times$ reduction in wall-clock time to the target metric.

Systems insight: Data echoing can trade sample diversity for accelerator utilization when the input pipeline is the bottleneck. The useful echo factor must be measured for the workload because it depends on the batch size, insertion point, augmentation, shuffling, and target metric.

Echoing can introduce correlations when the same example appears within or across nearby batches, so implementations must preserve adequate shuffling and validate the target metric. Choi et al. (2019) did not observe a negative batch-normalization interaction in their ResNet-50 or Single Shot MultiBox Detector (SSD) experiments, but that result does not guarantee the same behavior for every workload.

The engineering patterns in this section make data selection deployable. Proxy selection can reduce scoring cost when its ranking transfers to the target model. Sharded formats and shuffle buffers reconcile randomized sampling with sequential storage access, while data echoing can recover idle accelerator time when input service is the bottleneck. Each pattern changes a different part of the pipeline, and each still requires workload-specific measurement of quality and end-to-end time.

Engineering patterns solve the *how*, but a more fundamental question remains: whether to invest in data selection at all. A deduplication pipeline that costs \$50K to build but saves \$10K per training run requires a cost model to justify.

9.8 Cost Modeling

That investment decision demands cost modeling and quantitative answers. Practitioners must determine whether to label 10,000 more samples or buy more GPU hours, when active learning pays for itself, and what ROI a deduplication infrastructure investment delivers.

9.8.1 Quantifying data costs and ROI

Answering these questions requires understanding what training data costs. The total cost of data encompasses the full lifecycle of data acquisition, preparation, and utilization, extending well beyond storage fees. Table 9.14 breaks down the four cost components. Labeling is usually the component data selection can change most directly, while storage and processing tend to scale more mechanically with retained volume and training passes.

$$C_{\text{total}} = C_{\text{acquire}} + C_{\text{label}} + C_{\text{store}} + C_{\text{process}}$$

Table 9.14: Total Cost of Training Data: The four cost components span the full data lifecycle. Here D is the total sample count, D_{labeled} is the labeled subset, $D_{\text{vol,store}}$ is stored byte volume, T_{months} is retention time, N_{epochs} is the number of training passes, and O_{sample} is per-sample training work. The ranges are scenario anchors rather than universal market prices.

Component	Formula	Illustrative Range
C_{acquire}	$D \times c_{\text{sample}}$	\$0.001–\$10/sample (web scrape vs. licensed)
C_{label}	$D_{\text{labeled}} \times c_{\text{label}}$	\$0.01–\$200/sample (crowd vs. expert)
C_{store}	$D_{\text{vol,store}} \times c_{\text{storage}} \times T_{\text{months}}$	\$0.02–\$0.10/GB/month
C_{process}	$D \times N_{\text{epochs}} \times O_{\text{sample}} \times c_{\text{FLOP}}$	Proportional to training FLOPs

The interplay of these four cost terms becomes concrete when evaluating an ImageNet-scale vision model training run.



Napkin Math 9.5: Cost breakdown: ImageNet-scale training

Table 9.15 itemizes one illustrative acquisition, labeling, storage, and compute scenario at ImageNet scale. The resulting ratio follows from the stated licensing, annotation, storage, runtime, and compute-cost assumptions; it is not a typical ratio for every supervised workload.

Table 9.15: Illustrative ImageNet-Scale Cost Breakdown: Itemized acquisition, labeling, storage, and compute costs under the stated scenario assumptions. Each amount follows from the rate shown beside it, using the book’s published crowd-labeling, object-storage, and cloud-GPU rates. Data costs dominate this particular calculation because licensing and annotation are included while the training run itself is short.

Cost Component	Calculation	Amount
Raw data (1.2M images)	Licensed dataset, flat fee	\$50,000
Labels (crowd annotation)	$1.2\text{M} \times \$0.05/\text{label}$	\$60,000
Storage (cloud object store)	$150\text{ GB} \times \$0.02/\text{GB}/\text{month} \times 12\text{ months}$	\$36
Training campaign (30 runs, each 100 epochs in 24 h on 8 A100s)	$5,760\text{ GPU-hours} \times \$4/\text{GPU-hour}$	\$23,040
Total		\$133,076
Data vs. Compute ratio		82.7% data, 17.3% compute

Systems insight: Acquisition and labeling can dwarf compute cost in supervised vision workloads, as they do under these assumptions. The compute figure prices the whole training campaign, not one run: a published model is the survivor of sweeps, restarts, and ablations, and costing a single run would understate compute roughly thirtyfold. Even so, acquisition and labeling remain the larger share. The run count is an assumption like any other, so a complete ROI calculation must state it and price both categories rather than assuming which one dominates.

9.8.2 ROI framework for data selection techniques

Understanding total costs enables rational decisions about which efficiency techniques merit investment. Every technique carries both a cost (implementation effort, compute overhead) and a

benefit (reduced data requirements, faster training). Comparing these trade-offs requires a common framework: Return on Investment (ROI).

$$\text{ROI} = \frac{\text{Savings} - \text{Investment}}{\text{Investment}} \times 100\%$$

The challenge lies in quantifying both sides accurately. Table 9.16 summarizes the cost and benefit categories to measure; their ranking depends on the corpus, task, and existing infrastructure.

Table 9.16: ROI Profiles for Data Selection Techniques: The table identifies costs and possible savings. No technique guarantees positive ROI; each must satisfy the selection inequality under measured workload conditions.

Technique	Investment (Cost)	Savings (Benefit)
Deduplication	One-time compute for hashing + infrastructure	Reduced storage and repeated-sample processing
Coreset Selection	Proxy model training + selection compute	Fewer retained samples if target quality and coverage hold
Active Learning	Inference on unlabeled pool + human-in-the-loop latency	Lower labeling demand when query selection transfers
Data Augmentation	CPU/GPU cycles for transforms	Effective dataset size increase without new data acquisition

9.8.3 Break-even analysis

ROI calculations assume that techniques deliver their promised benefits, but actual outcomes vary. For any technique, there exists a **break-even point** where investment equals savings. Below this threshold, the technique costs more than it saves; above it, the technique generates value. Identifying this threshold determines whether a technique makes sense for a given project.

Suppose labeling costs \$10/sample, active learning starts from 1,000 labeled samples (\$10,000), issues 100 queries per round at \$50 inference cost, and the random-labeling baseline requires 5,000 samples for target accuracy. If active learning reaches target accuracy with only 2,000 labeled samples, the ROI follows from comparing labeling and compute costs.

$$\text{Random labeling cost} = 5,000 \times \$10/\text{sample} = \$50,000$$

$$\text{Active learning cost} = 2,000 \times \$10/\text{sample} + 10 \text{ rounds} \times \$50 = \$20,500$$

$$\text{ROI} = \frac{\$50,000 - \$20,500}{\$20,500} = 143.9\%$$

Break-even occurs when avoided-label cost equals selection overhead. A 20 percent labeling reduction may still yield negative ROI if scoring and retraining are expensive.

9.8.4 Amortization across training runs

Break-even analysis captures a snapshot in time, but many data selection investments span multiple projects. Techniques with high upfront costs yield significant returns when their benefits compound across repeated training runs. **Amortized return on investment (ROI)** accounts for this temporal dimension, as table 9.17 and table 9.18 illustrate for a deduplication pipeline:

$$\text{Amortized ROI} = \frac{N_{\text{runs}} \times \text{Per-Run Savings} - \text{One-Time Investment}}{\text{One-Time Investment}} \times 100\%$$

Table 9.17: Deduplication Infrastructure Cost Components: The one-time investment covers engineering effort and initial compute; the per-run savings accrue with every subsequent training run on the deduplicated data.

Component	Cost
Build deduplication pipeline	\$50,000 (engineering time)
Compute MinHash signatures (one-time)	\$5,000
Per-run savings	\$10,000/run

The ROI pattern in table 9.18 reveals which circumstances favor infrastructure investment. Three conditions tend to strengthen the return from data-selection investments:

- **Repeated training runs:** Hyperparameter search, model iterations, and scheduled retraining reuse the same selection infrastructure many times.
- **Shared datasets:** A cleaned or deduplicated corpus can support multiple teams or model architectures.
- **Broadly reusable techniques:** Methods such as deduplication transfer across models, whereas task-specific coresets may not.

Table 9.18: Amortized ROI over Multiple Training Runs: Under this scenario’s costs and savings, a deduplication pipeline that loses money on its first use becomes profitable when reused across enough runs. Infrastructure investments should be evaluated over their expected lifetime rather than from one use.

Number of Runs	Amortized ROI
1 run	-81.8% (net loss)
5 runs	-9.1% (near break-even)
10 runs	+81.8% (positive)
50 runs	+809.1% (highly profitable)

Deduplication can be a high-transfer investment because every model using the corpus inherits the storage and repeated-content changes, although quality effects can still differ by task. Task-specific coresets may transfer less reliably across architectures, limiting their amortization potential. For one-off runs, simple techniques such as random sampling or basic augmentation may yield better ROI than methods requiring substantial infrastructure.

The investment decision therefore reduces to a practical split between high-reuse, data-limited workflows and one-off, already-curated ones.

These ROI calculations all assume a single machine. Production ML training distributes data across many workers, introducing coordination overhead that can erode or amplify those returns: a coreset algorithm designed for a single GPU may behave differently once its dataset is sharded across hundreds of workers.

9.9 Distributed Selection

The earlier sections in this chapter used a centralized view in which one process can see the full dataset, compute global statistics, and coordinate selection. That abstraction lets coreset methods rank all samples, curricula establish a shared ordering, and active learning compare a common candidate pool. Distributed training breaks the abstraction when data is sharded and scores are refreshed at different model states. Two difficult problems arise: approximating a global selection from local views and keeping training-time rankings comparable as the model evolves.

Whether distributed selection stays faithful to the global dataset or collapses into local shard heuristics depends on how much cross-worker coordination each technique requires against the bandwidth available to sustain it. The selection problem therefore has to be evaluated at the boundary where statistical value meets sharding and locality.

These are independent requirements. A coreset can remain representative yet cost too much to coordinate, while a cheap shard-local method can scale but systematically miss rare groups visible only in the global corpus. A distributed design therefore needs two acceptance tests: compare selected-data quality with a centralized or otherwise auditable reference, and compare selection overhead with the end-to-end training time it saves. Passing only one test is insufficient.

9.9.1 Strategies for distributed selection

In a basic data-parallel layout, each worker processes a distinct shard and model synchronization is handled separately. Data selection can add dependencies across those shards (table 9.19).

The selection dependencies admit several architectural solutions, each navigating a different point in the consistency-scalability trade-off space. The most straightforward approach centralizes selection while distributing training. A coordinator node performs selection on the full dataset, then distributes selected indices to workers. This preserves selection quality but introduces a single bottleneck:

```
Coordinator: score_all_samples() → selected_indices
Broadcast: selected_indices → all workers
Workers: train on subset(local_shard, selected_indices)
```

Table 9.19: Selection Dependencies in Distributed Training: The table starts from a common centralized formulation of each technique and identifies the coordination problem introduced by sharding. Distributed variants may relax rather than reproduce the centralized result exactly.

Technique	Single-Node Assumption	Distributed Challenge
Coreset Selection	Global view of dataset	Each worker sees only its shard
Active Learning	Centralized uncertainty scoring	Scoring requires model synchronization
Curriculum Learning	Global difficulty ordering	Workers may have different “hardest” samples
Deduplication	Hash table fits in memory	Distributed hash tables add latency

The semantics remain clean, but the coordinator becomes a single point of failure and a possible bandwidth bottleneck. Whether that overhead is acceptable depends on the selection payload, refresh rate, network path, and cluster size.

Hierarchical selection addresses this scalability limitation by distributing the selection computation itself. Each worker performs local selection on its shard, then a coordinator merges results:

```
Workers: local_selected = select_top_k(local_shard)
Coordinator: global_selected = merge_and_rerank(all local_selected)
Broadcast: final_indices → all workers
```

Shard-local selection reduces coordinator load but introduces a quality trade-off: local quotas and score distributions may not preserve the global ranking or rare groups spread unevenly across shards. The merge step therefore needs normalization, coverage constraints, or a second global pass.

When even hierarchical approaches prove too expensive, approximate global selection offers a fallback. These methods trade exactness for scalability through distributed approximate algorithms. Distributed MinHash enables deduplication by having each worker compute MinHash signatures independently; signatures are then aggregated to find near-duplicates across shards without requiring any single node to see all the data. Similarly, distributed uncertainty sampling allows workers to compute local uncertainty scores, with a global threshold determined by score distribution statistics rather than exact ranking.

9.9.2 Consistency challenges in active learning

The approximate selection strategies assume static selection criteria, but active learning introduces an additional complication: the model changes during selection. Consider what happens when Worker A scores samples using the model at step t while Worker B simultaneously updates the model to step $t + 1$. Worker A’s scores are now stale and may select samples that the updated model would rank differently.

The scoring schedule therefore defines the meaning of a query, not merely its cost. A reproducible system versions the model checkpoint, candidate-pool snapshot, scoring rule, and selected indices together. It can then measure how much ranking agreement decays with checkpoint age and choose a refresh interval from that evidence rather than from an arbitrary number of steps.

Several strategies mitigate this staleness problem, each with distinct overhead characteristics:

- **Synchronous scoring:** All workers pause training and score simultaneously, guaranteeing consistency but at substantial cost in GPU utilization.
- **Periodic score refresh:** Workers re-score every k epochs rather than every batch, trading freshness for reduced overhead.
- **Checkpoint-robust selection:** The system selects samples that exhibit high uncertainty under multiple model checkpoints, keeping selection decisions valid as the model evolves.

An illustrative 8-node GPU-cluster scenario demonstrates how these refresh strategies interact with storage bandwidth and compute efficiency.

Example 9.6: Distributed coreset selection

Scenario: An 8-node A100 GPU cluster selects a 10 percent coreset from ImageNet (1.3M images) across distributed workers.

Diagnosis: Centralized coreset scoring on a single node creates memory and compute bottlenecks. Distributed pipeline staging (shard-local embedding → parallel deduplication → local EL2N scoring → centralized top- k merge) processes dataset selection in 67 minutes.

Systems lesson: Distributed coreset selection can remove single-node memory as the scoring limit, as illustrated in figure 9.11. Sharding embedding and scoring spreads work across the cluster, while the centralized merge remains a coordination cost that must satisfy the selection inequality.




Figure 9.11: Distributed Coreset Selection Architecture: Parallel shard-local scoring reduces the work assigned to a centralized rank-merging coordinator. Workers send local EL2N scores upward; the coordinator merges them and broadcasts final coreset indices. The design is beneficial only when this coordination overhead remains below the training time saved by the coreset.

Positive ROI can erode quickly when workers coordinate frequently during training. Distributed selection introduces a coordination tax whose size depends on what is synchronized and how often. That tax must remain smaller than the training time saved; if it approaches the limit, simplify the strategy or increase the refresh interval.

A further constraint arises from cluster network topology. Gradient synchronization may use accelerator collectives over dedicated links, while embedding vectors, score arrays, and coreset indices may follow host or storage paths instead. The actual route depends on system architecture. Selection traffic can therefore expose CPU, network interface card, storage, or collective-network bottlenecks that the training profile alone does not reveal.

Real ML systems combine data selection with model-level efficiency, machine-level throughput, and distributed training simultaneously. These optimizations interact in ways that can amplify or undermine each other, and understanding these interactions is essential for designing efficient end-to-end pipelines.

9.10 Cross-Layer Interactions

Data selection does not exist in isolation. A coreset-trained model may later be simplified for deployment. A curriculum-learning pipeline will run on specialized accelerators. An actively-learned dataset will feed into distributed training. These cross-layer interactions can amplify gains or introduce unexpected conflicts. Understanding these interactions helps practitioners design end-to-end efficient systems rather than optimizing components independently.

9.10.1 Model-level efficiency

Model-level efficiency reduces the size or arithmetic cost of the trained model through techniques such as pruning, lower-precision representation, and distillation. The training corpus can affect the learned representation and therefore the outcome of later compression, but a smaller or curated dataset does not guarantee a more compressible model. Chapter 10 and chapter 11 return to the hardware consequences of model simplification.

Systems Perspective 9.2: The sparsity latency trap

Context: The compressibility advantage from data selection is counted in arithmetic operations, but model compression research repeatedly shows that pruning reduces those counts more easily than it reduces wall-clock latency.

Failure mode: FLOPs can decrease dramatically while inference latency stays flat or increases. Dense matrix multiplication hardware rewards regular layout and reuse, while sparse matrices require irregular memory access, metadata checks, and address jumps. Unless the sparsity pattern matches the execution substrate, the overhead of managing sparsity can outweigh the reduction in arithmetic.

Systems insight: FLOPs are *not* latency. A 99 percent reduction in operations can yield a 0 percent reduction in time if the remaining operations are memory bound or cache-inefficient. Optimization must target the hardware's binding bottleneck rather than an abstract metric alone (Hoefler et al. 2021).

The mechanism relates to how models encode information. A model trained on repetitive data can learn redundant features that pruning later removes. The training compute required to learn those features was wasted, only to be discarded during compression. By contrast, a model trained on diverse, informative samples may learn compact, nonredundant representations from the start, making subsequent compression easier to evaluate and sometimes easier to apply. Treat this as an engineering hypothesis to measure after curation rather than a guaranteed property of every selected dataset. Data selection and model-level efficiency are complementary tools, but their interaction must be evaluated jointly by comparing compression quality and deployment performance after curation rather than assuming an ordering advantage.

9.10.2 Machine-level throughput

While model-level efficiency affects what work remains after training, machine-level throughput determines how efficiently training itself proceeds. Specialized accelerators, kernel optimization, and parallel execution increase effective throughput. Data selection affects which machine bottlenecks dominate, and this relationship is more nuanced than simple speedup calculations suggest, as table 9.20 illustrates.

Table 9.20: Diagnosing Bottlenecks after Data Selection: Dataset and access-pattern changes can expose a different bottleneck, but dataset size alone does not determine the limiting resource. Re-profile the pipeline and match the response to the measured condition.

Observed Condition	Possible Bottleneck	Candidate Response
Input rate below accelerator demand	Storage, decode, or host link	Sharding, prefetching, parallel decode
Low-intensity training kernels	Accelerator memory bandwidth	Fusion, reduced traffic, lower precision
Dynamic scoring dominates iteration	Selection compute or index I/O	Proxy models, cached scores or embeddings

Data selection can therefore shift the system from one bottleneck regime to another. A technique that reduces dataset size by 80 percent may expose input-pipeline latency, selection compute, or true GPU compute as the next bottleneck, requiring different optimizations in each case. Before applying aggressive data reduction, profile the system to understand which bottleneck is being targeted.

9.10.3 Distributed training

The hardware bottleneck analysis in section 9.10.2 assumes single-machine training. The interactions become more complex when scaling to multiple machines, because data selection affects different parallelism strategies in distinct ways.

With worker count, batch size, and epoch count fixed, fewer retained samples mean fewer gradient synchronizations and less communication over the full run. Step-based schedules or compensating epochs may remove that saving. Smaller per-worker shards can also improve locality, although access order, record size, cache capacity, and sharding determine the actual I/O effect.

The benefits must be weighed against the distributed selection challenges discussed in section 9.9. A technique that works well on a single GPU may incur prohibitive coordination overhead across 1,000 workers, negating its efficiency gains.

9.10.4 The optimization stack

The earlier subsections in this section examined pairwise interactions, but production systems apply all these optimizations together. Trace the full optimization stack in figure 9.12, from data to deployment: each stage in this pipeline amplifies or attenuates the effects of others.




Figure 9.12: The Optimization Stack: Data artifacts flow through selection and training before the resulting model enters compression and deployment. Upstream changes propagate quality and cost effects downstream, but reducing training data does not automatically simplify the model or deployment hardware.

The pipeline in figure 9.12 reveals why data selection occupies a strategic position at the head of the optimization stack. With the training schedule held fixed, reducing the dataset by 50 percent can halve sample-level training work. It does not by itself shrink the trained model, simplify compression, or relax serving hardware requirements. Each downstream stage inherits any efficiency gains or quality losses from upstream decisions, and poor selection can force compensation through longer training or less aggressive model simplification.

Quantifying this multiplicative effect requires determining whether a 50 percent dataset reduction delivers 50 percent compute savings, or whether it has inadvertently degraded model quality in ways that surface only in production. Answering this question requires a rigorous measurement framework: metrics that capture both the efficiency gains and the quality costs of data selection decisions.

9.11 Measurement Framework

The cross-layer stack makes a measurement framework unavoidable: every data selection technique claims to improve efficiency, but only rigorous measurement separates real savings from shifted costs or hidden quality loss. The core metrics in section 9.11.1 tie sample reduction to accuracy, cost, and deployment coverage so that a smaller dataset is judged by what it preserves, rather than by what it removes alone.

9.11.1 Core metrics

The core metrics connect sample reduction to model quality instead of reporting accuracy alone. They measure learning per sample, performance-per-data, and compression at a target accuracy.

Performance-per-data

The most direct metric, **performance-per-data (PPD)**, measures accuracy gain per sample:

$$\text{PPD}(n) = \frac{\text{Accuracy}(n) - \text{Accuracy}(0)}{n}$$

where n is the number of training samples. A higher PPD indicates greater average improvement per sample over the stated interval. Its shape must be measured because diminishing returns and their onset are workload-dependent.

Area under the learning curve

Rather than comparing at a single point, the **area under the learning curve** (AULC) integrates performance across all dataset sizes:

$$\text{AULC} = \int_0^D \text{Accuracy}(n) dn$$

where n is the dataset size and D is the total dataset size.

For the same metric and integration range, a higher AULC means the strategy achieves stronger performance with fewer samples. Comparisons must use the same upper limit D or normalize the integral.

Data compression ratio

For coreset methods, the **data compression ratio** (DCR) measures how much data reduction is achieved at a target accuracy:

$$\text{DCR} = \frac{D_{\text{full}}}{D_{\text{coreset}}} \text{ at Accuracy}_{\text{target}}$$

A DCR of $5\times$ means the coreset achieves target accuracy with 20 percent of the data.

9.11.2 The compute-optimal frontier

While the core metrics in section 9.11.1 measure individual techniques, a higher-level diagnostic is also needed to test whether the overall training strategy is data-limited or compute-limited. Neural scaling-law experiments (Kaplan et al. 2020; Hoffmann et al. 2022) fit power-law relationships over particular model families, datasets, and compute ranges. Controlled sweeps can use those fitted relationships to compare model and data allocations, but no single ratio diagnoses every training run.

The Chinchilla study³³ (Hoffmann et al. 2022) fit a compute-optimal balance between model size and training data within its experimental regime. Allocations on either side of that fitted balance used the study’s compute budget less effectively.

The optimal balance defines a **compute-optimal frontier**: the best achievable performance at each compute budget when data and model size are properly balanced. Figure 9.13 sketches this diagnostic as a conceptual frontier rather than a fitted Chinchilla result.




Figure 9.13: Illustrative Compute-Optimal Frontier: The green curve is a conceptual frontier over normalized compute, not a fitted result from Chinchilla. Points below it illustrate allocations that a controlled model-size and token-count sweep might identify as data- or compute-limited.

³³ | **Chinchilla:** Chinchilla is a 70-billion-parameter language model trained on 1.4 trillion tokens. In the experiments reported by Hoffmann et al. (2022), it outperformed the larger GPT-3 on most evaluated tasks under a similar training-compute budget. The fitted compute-optimal allocation scaled parameters and tokens at roughly equal rates within the study’s experimental regime; it is not a universal prescription for every architecture, dataset, or training recipe.

Against a frontier fitted from controlled sweeps, a point’s location can suggest the next experiment. A data-limited allocation motivates tests of corpus quality, coverage, or selection; a compute-limited

allocation motivates tests of effective throughput, duration, or model size. The conceptual points in figure 9.13 do not diagnose a real run by themselves, and points near a fitted frontier remain specific to the model family, data, objective, and compute range used to estimate it.

Interpreting the Chinchilla result

Within the fitted Chinchilla regime, compute-optimal model parameters and training tokens grew at roughly equal rates.³⁴ Combining that empirical fit with the dense-transformer compute approximation, we derive $D_{\text{opt}} \propto \sqrt{C}$, so doubling compute corresponds to about $\sqrt{2} - 1$, or 41.4 percent, more tokens under those assumptions. The commonly quoted 20 tokens per parameter is a useful reference point from that study, not a universal optimum.

Applying the diagnostic

If a training run underperforms expectations, extending one run can reveal whether additional optimization steps still help, but a plateau does not identify data starvation by itself. Learning-rate schedules, optimization limits, model capacity, and data quality can all produce a plateau. A reliable diagnosis compares controlled runs that vary model size, token count, and compute allocation while holding the evaluation protocol fixed.

In figure 9.14 the two curves diverge: a data-efficient selection strategy (blue) reaches the performance plateau with fewer samples than random sampling (gray). The horizontal gap represents that sample reduction and can translate into compute savings when the per-sample work and training schedule remain fixed; the vertical gap marked by the red arrow represents the performance gained at a fixed dataset size.




Figure 9.14: Illustrative Diminishing Returns of Data: The synthetic data-efficient curve (blue) reaches the same plateau with fewer samples than random sampling (gray). The vertical red arrow marks the modeled performance gap at a fixed dataset size; measured curves need not have this shape.

For practitioners, the metrics in this section answer a practical question: at what point to stop collecting data and start curating it, and when adding more samples wastes compute rather than improving accuracy. Knowing that a strategy is efficient, however, is not the same as confirming that a curated dataset preserved model quality.

Data selection techniques implicitly rank sample value, and validating that a curated dataset preserves model quality requires systematic benchmarking across three dimensions. Coverage metrics validate that coreset selection preserved representation across classes and demographic groups. Distribution alignment metrics (such as KL divergence and population stability index, which section B.2.2 defines) detect whether the curated training set drifted from the deployment distribution. Label quality metrics (inter-annotator agreement, confident learning) validate that active learning did not introduce systematic labeling errors. A 50 percent dataset reduction is only valuable if benchmarking confirms the model maintains target accuracy, calibration, and robustness. Chapter 12 later generalizes these ideas into broader model-and-data evaluation protocols; here,

³⁴ **Tokens per parameter:**

The ratio D/P compares training tokens with model parameters; GPT-3 used approximately 1.7 tokens per parameter (Brown et al. 2020), while Llama 2 70B used approximately 28.6 (Touvron, Martin, et al. 2023). These descriptive ratios do not establish whether either run is compute-optimal under a different architecture, corpus, objective, or hardware budget. Controlled model-size and token-count sweeps are required for that diagnosis.

the point is narrower: a curated dataset must be validated against the task distribution rather than against its reduction ratio alone.

The validation target can itself be unreliable. A widely cited replication study shows how benchmark accuracy can overstate what a model has learned, which is why the evaluation set deserves the same scrutiny as the curated training set.

 Example 9.7: The benchmark replication gap

Scenario: Classifiers evaluated on newly collected test sets experienced 3 to 15 percentage-point accuracy drops on CIFAR-10 and 11 to 14 points on ImageNet (Recht et al. 2019).

Diagnosis: The replicated collection produced slightly harder images than the original test sets. The study found that benchmark adaptivity did not explain the gap, even though repeated use of a fixed test set remains a general evaluation risk.

Systems lesson: Evaluation-set construction affects measured generalization. Data-selection pipelines should validate curated subsets on independently collected and deployment-relevant test sets rather than trusting one long-reused benchmark.

The lesson carries directly to data selection: an efficiency gain measured against a single benchmark can evaporate on newly collected data, so a curated subset must hold up across multiple held-out distributions before its reduction ratio is trusted.

 Lighthouse 9.3: Lighthouse data selection

Data selection principles apply to all five lighthouse models, though the dominant constraint determines the technique. Table 9.21 maps each workload to a priority: for example, compute-bound ResNet-50 benefits from coreset selection, memory-bound GPT-2/Llama from deduplication, and the tiny KWS model from augmentation and synthesis.

Table 9.21: Data Selection Priorities by Lighthouse: Primary bottleneck and candidate data-selection priority for each lighthouse workload. The pairing suggests a hypothesis to test; the dominant resource constraint alone does not guarantee that a technique pays off.

Lighthouse	Primary Bottleneck	Data Selection Priority
ResNet-50	Compute	Coreset selection directly reduces training FLOPs
GPT-2/Llama	Memory bandwidth	Deduplication reduces corpus size; curriculum learning improves token efficiency
MobileNetV2	Latency/Power	Validate augmentation against the target model and deployment invariances
DLRM	Memory capacity	Interaction deduplication and embedding pruning reduce table size
Keyword Spotting	Extreme constraints	Augmentation and synthesis create datasets from minimal seeds

Across these models, data selection is not a single technique but a systems optimization tailored to whichever resource is most constrained.

The measurement tools and lighthouse examples demonstrate what data selection can achieve when applied correctly. The techniques, however, involve counterintuitive trade-offs, and practitioners frequently fall into predictable traps.

9.12 Fallacies and Pitfalls


Data selection involves counterintuitive diminishing returns that contradict the “more is better” intuition from traditional machine learning. The following errors fall into three groups: conceptual fallacies about what data selection can achieve, implementation pitfalls that arise when correct strategies meet engineering realities, and transfer errors that occur when benchmark results are applied uncritically to new domains.

Fallacy: *Data is the new oil, so more is always better.*

More data does not imply proportional accuracy gains because learning curves commonly show diminishing returns. The numerical comparison in this callout is illustrative rather than measured: it assumes a move from 1M to 10M samples for only 4 percentage points of accuracy gain. Table 9.1 likewise uses scenario growth rates. Teams should measure the learning curve and marginal compute cost for their own data rather than infer either from these example values.

Pitfall: *Replacing real-data validation with synthetic-only training data.*

Engineers assume generative models can replace data collection with inexhaustible generated examples at marginal cost. Synthetic-only training can fail through two different mechanisms. First, section 9.5.3 and table 9.9 show the domain-gap problem: generated data can diverge from the real deployment distribution, causing the learned decision boundary in figure 9.8 to misclassify real-world inputs. Second, recursive training on model-generated data can cause model collapse (Shumailov et al. 2024): in this illustrative scenario, accuracy degrades from 95 percent to 78 percent after five generations, a 17 percentage-point drop. The illustrative 50–80 percent synthetic range is workload-dependent and must be validated against real deployment data.



Fallacy: *Data selection is merely data cleaning.*

Engineers can conflate data quality, which includes removing errors, with data value under a particular task and training procedure. Figure 9.4 illustrates one uncertainty-based strategy, while section 9.2.2 also covers geometric objectives. The 1.8× ICR difference is a worked scenario rather than a general EL2N or GraNd result. Cleaning addresses data defects; selection evaluates the learning value and cost of retained examples.

Pitfall: *Treating data selection as a budget-only tactic.*

Practitioners may view data selection as relevant only for TinyML or budget-limited teams. It applies wherever high-quality examples, not raw volume, constrain learning. A 10 percent efficiency gain on a \$100M training run saves \$10M under the scenario assumptions. The data wall (figure 9.1) can become especially visible at large scale, when compute capacity outruns accessible, legally usable, target-relevant data. Section 9.8.4 shows why reuse can make selection infrastructure economical even for well-funded teams.

Conceptual misunderstandings often lead to flawed strategies. Equally damaging are the implementation pitfalls that arise when correct strategies meet messy engineering realities.

Fallacy: *Selection overhead is too small to change training economics.*

Take a scenario in which the full run takes 8 hours and training on the coreset takes 2-hour, so selection has to buy back its cost out of the hours saved. A sophisticated coreset algorithm requiring 10 hours to do that selection spends 5× the training run it enables and more than the full run it was meant to avoid, yielding negative ROI. Equation 9.2 defines the boundary. Use lightweight proxy models or cached embeddings when they preserve selection quality. In the illustrative comparison, proxy-based EL2N scoring completes in 30 minutes, or 6.2 percent of that full-run baseline, and satisfies the inequality.

Pitfall: *Pruning rare classes into oblivion.*

Aggressive selection without group constraints can remove rare classes because an average-loss objective gives them little weight. In a 1M dataset with 0.1 percent rare-class samples (1,000 examples), an unstratified 10 percent subset retains only 100 such examples in expectation, below the scenario's 150-sample target. **Importance sampling** can reweight rare samples, but large weights can inflate gradient variance. Section 9.2.2 therefore recommends minimum representation by class or deployment-relevant slice before optimizing the remaining selection budget.

Fallacy: *Deduplicating training data is enough to make evaluation reliable.*

Some standard language-modeling datasets have measurable train-test overlap; a deduplication study reports overlap affecting over 4 percent of validation examples in evaluated datasets (Lee et al. 2022). Deduplicating only training data is therefore insufficient: evaluation sets must also be checked against the training corpus. This extends the threshold-calibration principle in section 9.2.3

across dataset splits. Teams should validate the effect because near-duplicate thresholds and domain distributions matter.

Pitfall: *Active learning without considering annotation latency.*

Active-learning analyses often abstract away oracle response time. In real annotation workflows, turnaround can affect the usefulness of a query batch. The 14-day latency and batch sizes in this callout are illustrative; the correct batch size depends on measured annotation capacity, model-update cadence, and diversity requirements. Active learning ROI therefore depends on both label cost and turnaround time.

Fallacy: *If a technique works on ImageNet, it will work on my dataset.*

Data-selection effectiveness depends on dataset redundancy, representation, model, and target metric. Results from CIFAR-10 or ImageNet do not establish a safe pruning ratio for medical, satellite, or other domain-specific data, and those datasets cannot be assumed to have near-zero redundancy either. Start with a measured baseline, vary the retained fraction, and validate overall and slice-level quality before aggressive reduction.

Pitfall: *Optimizing data selection metrics instead of deployment metrics.*

The illustrative failure scenario uses a 10 percent coreset that performs well on majority classes but poorly on rare subgroups. Section 9.11.1 defines efficiency metrics, while section 9.11 requires stratified evaluation. Selection must preserve demographic groups, rare classes, and deployment failure modes even when doing so reduces average PPD.

9.13 Summary

These fallacies and pitfalls arise when practitioners treat data selection as a purely algorithmic exercise divorced from the systems context in which it operates. Smaller, curated datasets sometimes outperform massive ones because data selection is a *systems* problem rather than a purely *statistical* one. It addresses the first question in the optimization ordering by reducing work before it begins. Traditional machine learning frames the problem as the number of samples needed to achieve target accuracy; the systems perspective frames it as minimizing total cost across the entire pipeline.

This systems framing becomes actionable through the ICR metric, the selection inequality, and the cost modeling framework. The goal is to minimize total cost across compute, storage, labeling, energy, and time rather than maximize accuracy alone.

The three-stage optimization pipeline addresses different phases of this cost equation: static pruning removes redundancy before training through coreset selection and deduplication, dynamic selection prioritizes informative examples during training through curriculum and active learning, and synthetic generation creates data where none exists through augmentation, simulation, and distillation. Together, these strategies address the “data wall,” the structural asymmetry between rapidly growing compute capacity and slowly growing high-quality data.

Self-supervised learning occupies one end of the data-selection spectrum: pretraining shifts much of the learning signal into a reusable model, so downstream tasks can often use fewer task-specific labels. The economics improve when the pretraining cost is amortized across enough tasks.

Translating these techniques into production requires systems engineering: the selection inequality in equation 9.2 gates every technique, proxy models and shard-based data loaders reconcile selection algorithms with storage hardware, and data echoing can recover idle GPU cycles when input service is the bottleneck. The cost modeling framework (total data cost, ROI analysis, and break-even thresholds) provides the quantitative tools to evaluate which techniques merit investment for a given workload, while core metrics (PPD, AULC, DCR) and a fitted compute-optimal frontier help practitioners test whether training is data-limited or compute-limited.



Key Takeaways: Curate, do not accumulate

- **Selection optimizes system cost:** The goal is reduced total cost across the entire pipeline (compute, storage, labeling, energy), rather than “fewer samples for same accuracy” alone. The information-compute ratio quantifies learning gained per FLOP spent. Its

time benefit matches a throughput improvement only under the corresponding compute-bound assumptions.

- **Start with deduplication:** Exact deduplication is often the lowest-risk first technique because it removes repeated bytes without a learned scoring pass. Apply near-duplicate and semantic deduplication only after validating thresholds against downstream quality metrics.
- **The selection inequality gates every technique:** $T_{\text{selection}} + T_{\text{train}}(\text{subset}) < T_{\text{train}}(\text{full})$. Selection overhead must remain below the training time saved by the subset. Proxy models and cached embeddings can keep $T_{\text{selection}}$ low; expensive selection algorithms can consume all the savings they promise.
- **Dynamic selection adapts the data diet as the model learns:** Curriculum learning changes presentation order, while active learning changes which samples receive labels. Either can improve efficiency when its scoring rule, coverage, and refresh overhead are validated.
- **Self-supervised pretraining can amortize adaptation cost:** Pretraining once and fine-tuning many times spreads pretraining cost across downstream tasks. The worked scenario assumes a $100\times$ label reduction and a $20\times$ marginal-compute reduction; realized savings depend on transfer quality and reuse.
- **Synthetic data is a supplement, not a replacement:** The 50–80 percent synthetic mixture is illustrative and workload-dependent, not a universal optimum. Pure synthetic training risks model collapse and domain-gap degradation.
- **Validated workload reduction compounds downstream:** When selection safely eliminates examples or tokens without adding compensating steps, that work never reaches later model or machine optimizations. Dynamic scoring and synthesis add their own costs, so the combined saving must be measured rather than assumed.


The techniques explored throughout this chapter (deduplication, coreset selection, curriculum learning, active learning, and synthetic generation) provide a toolkit for addressing the data wall. When validated for the workload, they can reduce labeling, iteration, storage, or training costs while preserving the coverage needed for generalization.

The answer to the chapter’s opening question is that examples can differ substantially in learning value. Some benchmark datasets can shed half their examples with no measured loss under a validated selection method, while more aggressive pruning usually trades additional savings for some quality loss. Selection is the discipline of distinguishing informative, representative examples from redundant or harmful ones without removing rare cases that matter at deployment. Through D·A·M, this is data-algorithm co-design working upstream of everything else, because a sample excluded safely from training is a cost no downstream optimization has to absorb.

What’s Next: From data to algorithms

Data selection changes *what* the system learns from; even the best data cannot make an inefficient model run fast on constrained hardware. Chapter 10 shifts the focus to *how* the system represents what it learns, applying pruning, quantization, and knowledge distillation to reduce the computational cost of the model artifact itself.

Model Compression




- 10.1 Optimization Framework
- 10.2 Deployment Context
- 10.3 Structural Optimization
- 10.4 Quantization and Precision
- 10.5 Architectural Efficiency
- 10.6 Technique Selection
- 10.7 Optimization Strategies
- 10.8 Efficiency Measurement
- 10.9 Implementation Tools
- 10.10 Fallacies and Pitfalls
- 10.11 Summary

Purpose

Why do the models that win benchmarks rarely become the models that run in production?

Training produced a capable model, yet capability alone does not guarantee deployability. Cloud, Edge, Mobile, and TinyML each impose constraints that research benchmarks ignore, including memory budgets measured in megabytes rather than gigabytes, latency targets measured in milliseconds rather than seconds, and power envelopes measured in milliwatts rather than kilowatts. Research optimizes for accuracy on held-out test sets; production optimizes for accuracy per dollar, accuracy per watt, and accuracy per millisecond. Models that win benchmarks are often larger, slower, and more resource-intensive than production constraints permit. Bridging that gap requires a systematic discipline of compression that trades capabilities the deployment does not need for constraints it cannot violate. Many trained models carry more precision, more connections, or more capacity than a deployment context demands, and some of that surplus can be removed while preserving required behavior. Yet a smaller representation is not automatically a faster one. The hardware and software stack must exploit the new precision or structure; otherwise a nominal reduction can merely move the bottleneck while latency remains unchanged. Applied well, compression can substantially reduce model size, transforming a research artifact that runs only in a data center into a production asset for a phone, sensor, or microcontroller. The discipline is not simply about making models smaller but about making the right models possible for their physical environment. In D·A·M terms, compression enacts algorithm-machine co-design on the model itself, rewriting its mathematical structure to fit the physical constraints of the machine.





Learning Objectives

- Explain compression as algorithm-machine co-design that trades surplus capacity for memory, latency, and energy constraints
- Compare pruning, distillation, quantization, and architecture search by the resource constraint each relaxes
- Calculate parameter memory, precision, and sparsity reductions to estimate best-case compression gains
- Apply post-training, quantization-aware, and weight-only strategies under accuracy and hardware constraints
- Select structured pruning and operator choices that map to available accelerator kernels
- Design compression pipelines that order pruning, distillation, and quantization to preserve deployment accuracy
- Evaluate measured latency, energy, and accuracy on target hardware rather than relying on FLOP counts

10.1 Optimization Framework

A 7-billion parameter language model requires 14 GB merely to store its weights in FP16. The deployment target is a smartphone with 8 GB of RAM shared across the operating system, applications, and the model. *The math does not work.* No amount of clever engineering changes this arithmetic: 14 GB *cannot* fit in 8 GB. Yet users expect the model to run responsively, offline, without draining their battery in an hour. Every request has only a small time window in which to load data, run arithmetic, and return a result; the broader deployment gap also includes memory capacity, energy, and offline execution. That gap is not a minor inconvenience but a defining challenge of model compression.

Recall the silicon contract (principle 4), the performance bargain a model makes with its hardware. Compute throughput, memory bandwidth, memory capacity, and fixed overhead determine which resource binds first. During training, this contract is often negotiated upward. Researchers can select larger architectures, higher numerical precision, and deeper layers when a GPU cluster can afford those demands. Training also carries gradients, optimizer state, and numerical headroom that a forward-only deployment path may not need. In section 8.6.3, mixed precision improves training efficiency while preserving the numerical behavior needed to learn. Compression goes further for deployment by reducing precision to INT8 or below when the target execution path supports it. Deployment may reverse the training priorities: the production environment can be smaller, power-constrained, and latency-sensitive, even when the model was developed without those limitations. Where data selection optimized what the model learns from, compression optimizes what the trained model carries into that environment. **model compression** is the systematic process of renegotiating that contract for a new execution context, reducing memory footprint, computational cost, or energy consumption while preserving the behavior the application requires.

The scale of this renegotiation makes model optimization an engineering discipline, not a collection of ad hoc tricks. A 175 billion parameter model consumes over 350 GB in FP16 representation alone, while a microcontroller offers only 512 KB of SRAM. Bridging six orders of magnitude requires systematic methods with predictable trade-offs, not trial and error. Every optimization technique removes something from the model (redundant parameters, numerical precision, or architectural complexity), and we must understand exactly what is lost, what is preserved, and how these losses compose when techniques are combined.

Compression works along three complementary dimensions. Structural optimization removes redundancy from the model itself: **Pruning** eliminates low-impact parameters, knowledge distillation transfers behavior into a smaller architecture, and neural architecture search discovers designs for a specified objective. Precision optimization reduces the bit width of weights and activations; for example, FP32-to-INT8 conversion cuts the raw bytes per represented value by four. Supported low-precision matrix units can also accelerate arithmetic. Hardware-level optimization maps the result to the target processor through techniques such as fusion and supported sparsity. These

175B FP16 350 GB

Phone RAM 8 GB

MCU RAM 512 KB

log scale

Frontier weights dwarf phone and microcontroller memory; compression bridges the gap.

dimensions form an optimization stack: structural changes alter which operations exist, precision changes bytes per value, and hardware-level work determines whether those theoretical savings become supported execution paths. A practitioner might prune ResNet-50 filters, quantize the remainder to INT8, and fuse batch normalization into convolution. Benefits compound only when the transformations relieve different bottlenecks and the runtime supports the resulting graph. Each layer therefore needs its own measurement. Section 11.4.3 explains the accelerator mechanisms behind low-precision paths.

Concrete systems keep those trade-offs measurable: ResNet-50 and MobileNetV2 (our lighthouse models from section 6.1.1) for vision workloads, transformer-based language models for sequence tasks, DLRM for recommendation memory pressure (Naumov et al. 2019), and the depthwise-separable convolutional neural network (CNN) known as DS-CNN for TinyML keyword spotting (Y. Zhang et al. 2017). Reusing these models lets us compare techniques under consistent conditions, making the trade-offs between accuracy, latency, memory, and energy tangible rather than abstract.



Definition 10.1: Model compression

Model compression is a family of techniques that reduce a trained model's computational cost and memory footprint by eliminating redundant parameters (pruning), reducing numerical precision (quantization), or transferring learned behavior into a smaller architecture (distillation), while preserving as much predictive accuracy as possible.

1. **Significance:** Compression directly reduces the iron law's data-movement and compute terms. INT8 quantization of a 175-billion-parameter LLM cuts weight memory from 350 GB (FP16) to 175 GB, a 2× reduction in D_{vol} , while dedicated low-precision matrix units can increase compute throughput when kernels and layouts use the supported INT8 path. Unstructured pruning to 50 percent sparsity halves the nonzero count, but executed operations fall only when the runtime can skip those zeros through a supported sparse format or kernel.
2. **Distinction:** Unlike post-training compression methods such as pruning and quantization, neural architecture search discovers efficient architectures from scratch by exploring a design space. Here, NAS is treated as a related structural optimization technique: it changes the representation before training rather than compressing a finished model post hoc.
3. **Common pitfall:** Compression techniques do not compose without interference. Pruning changes weight distributions and operation patterns; quantization adds calibration and kernel constraints. Their combination can lose accuracy or fail to accelerate without joint validation.

The optimization stack moves from representation to numerics to execution. Deployment context determines which constraint binds first; structural methods change the computation, precision methods change the representation of each value, and architectural methods decide whether the compressed artifact actually maps to efficient hardware execution. Selection and composition follow from that constraint order rather than from a checklist of techniques.

The three dimensions form a natural hierarchy. Optimization first determines *what* computations the model should perform (representation), then *how precisely* to perform them (numerics), and finally *how efficiently* to execute them on physical hardware (implementation). Figure 10.1 shows this progression from software concerns toward hardware-level execution.

The top layer, efficient model representation, focuses on eliminating redundancy in the model structure. Techniques like pruning, knowledge distillation, and neural architecture search (NAS)¹ reduce the number of parameters or operations required, addressing memory footprint and computational complexity at the algorithmic level.

The middle layer, efficient numerics representation, optimizes how numerical values are stored and processed. **Quantization** and mixed-precision training reduce the bit-width of weights and activations (for example, from 32-bit floating point to 8-bit integers), enabling faster execution and lower memory usage on specialized hardware.

¹ | **Neural architecture search (NAS):** Zoph and Le (2016) at Google Brain used reinforcement learning to learn the architecture itself at a cost of 22,400 GPU-days (800 GPUs for 28 days), equivalent to 537,600 GPU-hours. Weight-sharing approaches such as Efficient Neural Architecture Search (ENAS) later reduced search cost by roughly 1,000× by sharing parameters across candidate architectures (Pham et al. 2018). Hardware-aware NAS and scaling methods then made the search output practical for deployable architecture families such as EfficientNet and MobileNetV3 (Tan and Le 2019; Howard et al. 2019).




Figure 10.1: Optimization Stack: Model optimization progresses from software-level representation (pruning, distillation, NAS) down through numerical precision (quantization) to hardware execution (kernel fusion, sparsity engines). Lowering abstraction moves optimizations closer to physical silicon, where representation decisions dictate achievable execution efficiency.

The bottom layer, efficient hardware implementation, asks whether the remaining operations map efficiently to the target processor. Operator fusion, supported sparsity, and hardware-aware scheduling align computation with memory hierarchies and vector or matrix units; profiling determines whether they improve utilization or throughput.

These dimensions are interdependent. Pruning reduces complexity but may require architectural changes for hardware efficiency. Quantization reduces precision but changes numerical behavior and the available execution paths. A useful strategy may combine techniques across layers when they address different binding constraints. For practitioners seeking immediate guidance, section 10.6.2 maps deployment constraints to candidate techniques. The intervening sections provide the technical foundation needed to evaluate those choices.

The relative importance of each dimension varies by deployment target. Cloud systems may tolerate larger models but demand throughput; mobile devices prioritize memory and energy; embedded systems face hard constraints on all resources simultaneously. Understanding these deployment contexts shapes which optimization dimensions to prioritize.

10.2 Deployment Context

The preceding optimization framework identifies three dimensions of compression, but which dimensions matter most depends entirely on where the model will run. A data center GPU with 80 GB of high-bandwidth memory (HBM) faces different binding constraints than a smartphone with shared RAM or a microcontroller with only a few hundred kilobytes of SRAM. Table 10.1 summarizes the key constraints across deployment environments.

Table 10.1: Deployment Constraints: Each deployment context imposes different optimization priorities.

Context	Memory	Latency	Power	Primary Goal
Cloud	tens of GB	100–500 ms	Flexible	Throughput, cost
Mobile/Edge	hundreds of MB to GB	5–100 ms	W-scale	Size, latency
TinyML	KB–MB	1–10 ms	mW	Size, energy

10.2.1 Deployment scenarios

Cloud inference often centers on throughput (requests/second/dollar), where supported quantization paths can increase serving density and operator fusion can reduce per-request latency (Choudhary et al. 2020; Dean et al. 2018). Mobile and edge deployments must fit device memory while meeting real-time targets. A camera app processing 30 fps has 33 ms per frame, so crossing that latency threshold can determine whether the feature operates in real time.

TinyML makes optimization a deployment requirement. A microcontroller with a few hundred kilobytes of RAM *cannot* run a 100 MB model regardless of accuracy. The model must fit below hardware limits or deployment is impossible (Banbury et al. 2020). Mobile devices offer more headroom, but an optimization that crosses a memory, latency, or thermal threshold can still determine whether a feature ships.

This deployment-time pressure is not new. AlexNet, the landmark model that won the 2012 ImageNet challenge, encountered a memory wall during training, and its architecture records the resulting two-GPU design choice.

Example 10.1: AlexNet’s two-GPU split (2012)

Context: In 2012, Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton aimed to train a 60-million-parameter vision model on ImageNet (Krizhevsky et al. 2012).

Mechanism: An NVIDIA GTX 580 provided only 3 GB of GDDR5 memory—insufficient to store the 240 MB model weights alongside FP32 activations (M_{act}) and optimizer states (M_{state}) for single-GPU training.

Impact: The model could not be trained on a single device at all, so the memory ceiling propagated upward into the network design itself rather than remaining a hardware procurement problem.

Fix: They partitioned the model across two 3 GB GTX 580 GPUs and limited cross-GPU connections so that the working set would fit within each device’s memory bound ($M_{\text{peak}} \leq 3 \text{ GB}$). The resulting two-tower architecture achieved a winning top-5 error rate of 15.3 percent, a 10.8 percentage point margin.

Systems lesson: Memory capacity bounds have constrained deep learning since its inception. Model deployment strategies (pruning, quantization) address the same physical limits that forced model splitting in AlexNet.

The same memory pressure that shaped AlexNet’s architecture appears in everyday product constraints when the model must run on commodity mobile hardware.

Example 10.2: An illustrative MobileNet win

Scenario: Deploying a real-time background blur video segmentation model on mid-tier Android smartphones at 30 FPS.

Diagnosis: An uncompressed FP32 MobileNetV3 model achieves only 8 FPS, exceeding device latency and thermal power limits. Quantizing weights to INT8 enables mobile NPU/DSP integer matrix execution.

Systems lesson: Here, INT8 cuts raw parameter payload by $4\times$, and a supported integer path raises inference from 8 FPS to 35 FPS. Speed and power still require device measurement.

The gap between model requirements and device capabilities explains why compression is not optional for resource-constrained deployment.

10.2.2 Balancing trade-offs

Table 10.2 quantifies this deployment gap using the lighthouse models from section 6.1.1: even MobileNetV2 at INT8 precision exceeds TinyML device memory by about $6.7\times$. This mismatch makes the **accuracy-efficiency trade-off** unavoidable. Increasing model capacity generally enhances predictive performance while increasing computational cost, resulting in slower, more resource-intensive inference. The improvements introduce challenges related to memory footprint, inference latency, power consumption, and training efficiency.

Systems Perspective 10.1: The compression-accuracy trade-off curve

The frontier is not uniform: structure-preserving techniques sit near the top, exchanging little or no accuracy for real speedups, while aggressive structural surgery sits at the bottom, where each additional gain extracts a steepening accuracy penalty.

The engineering decision is where to stop. Compression should halt at the “knee” of the curve, the point where the marginal loss in accuracy first exceeds the marginal gain in efficiency. Past that knee, the model degrades faster than it accelerates.

Table 10.2: The Deployment Gap: Model weight requirements compared against typical device capacities. DLRM represents recommendation-model embedding pressure (Naumov et al. 2019); DS-CNN represents the TinyML keyword-spotting case (Y. Zhang et al. 2017). Even MobileNetV2 quantized to INT8 exceeds the ~512 KB TinyML envelope by 6.7×, while the purpose-built DS-CNN keyword spotter fits after INT8 quantization. In this weight-only accounting, runtime weight memory equals artifact weight storage; activation, workspace, packaging, and metadata overheads are excluded. Numbers in parentheses show how many times the model exceeds device memory.

Model	Runtime Weight Memory	Artifact Weight Storage	Cloud (~107 GB)	Mobile (8 GB)	TinyML (~512 KB)
DLRM	100 GB	100 GB	ok	no (11.6×)	no (190734.9×)
GPT-2 XL	6 GB	6 GB	ok	ok	no (11444.1×)
ResNet-50	102.4 MB	102.4 MB	ok	ok	no (195.3×)
MobileNetV2	14 MB	14 MB	ok	ok	no (26.7×)
MobileNetV2 (INT8)	3.5 MB	3.5 MB	ok	ok	no (6.7×)
DS-CNN (KWS, INT8)	200 KB	200 KB	ok	ok	ok

This tension manifests differently across deployment contexts. Training requires computational resources that scale with model size; inference demands strict latency and power constraints in real-time applications. Understanding where each optimization technique falls on the compression-accuracy Pareto frontier is essential for informed technique selection.

Table 10.3 summarizes the key optimization techniques, their systems benefits, and their ML costs. These are empirical relationships—actual results depend on model architecture, task, and careful implementation.

Table 10.3: The Optimization Trade-Offs: Region 1 = Free Lunch, Region 2 = Efficient Trade, Region 3 = Danger Zone. Batch size affects training dynamics rather than model quality directly. These are representative engineering ranges synthesized from published work on inference runtimes, quantization, pruning, low-bit LLM quantization, and distillation (NVIDIA 2024c; Jacob et al. 2018; Han et al. 2015; Lin et al. 2024; Hinton et al. 2015); actual results vary with architecture, task, calibration data, and implementation quality.

Technique	Systems Gain	ML Cost	Typical Impact	Region
Operator Fusion	10–30% latency reduction	None	No accuracy loss	1
FP32 → BF16	2× memory, ~2× throughput	Minimal	< 0.1% accuracy drop	1
FP16 → INT8	2× memory, 2–4× throughput	Quantization error	0.5–1% accuracy drop	2
50% Pruning	~2× smaller model	Capacity loss	0.5–1% accuracy drop	2
Knowledge Distillation	2–10× smaller student	Capability ceiling	1–3% accuracy drop	2
4-bit Quantization	4× memory reduction	Significant error	2–5% accuracy drop	2–3
90% Pruning	~10× smaller model	Severe capacity loss	5–15% accuracy drop	3
↑ Batch Size (8×)	Higher throughput, better GPU util	Generalization gap	Requires LR scaling	—

Techniques that preserve model structure, such as fusion and precision reduction, tend to be “free” or inexpensive, while techniques that alter structure, such as pruning and distillation, extract more savings but require careful tuning. Each deployment context imposes a different binding constraint, including memory capacity on mobile devices, latency in real-time systems, and energy on battery-powered sensors. The optimization stack follows those constraints downward. Structural methods modify what computations occur, reducing the model’s parameter count and operation count to fit tighter memory and compute budgets. Precision techniques reduce how many bits represent each value, directly shrinking memory footprint and accelerating arithmetic. Architectural

approaches improve how efficiently the remaining operations execute on physical hardware, closing the gap between theoretical savings and measured performance.

Checkpoint 10.1: The efficiency frontier

Optimization is about trading one resource for another.

Trade-offs

- ❑ **Accuracy vs. cost:** compare a structure-preserving optimization and a structure-altering one by naming which resource each saves and which quality risk it introduces.
- ❑ **The “free lunch”:** can you point to where on the compression-accuracy frontier a technique buys efficiency at no accuracy cost, and explain why that region exists?
- ❑ **The knee:** explain why aggressive optimization should stop where the marginal accuracy loss exceeds the marginal efficiency gain.

10.3 Structural Optimization

Modern neural networks often carry more parameters than deployment requires.² Unused capacity still consumes memory, computation, and energy. Structural optimization addresses the first dimension of our framework, **efficient model representation**, by modifying *what* the model computes. Although additional capacity can aid optimization during training, parameters that do not improve the deployed model still impose deployment cost.

Every technique in this chapter follows the same engineering heuristic: the **conservation of complexity**. Compression rarely destroys cost outright. It relocates cost between the Data, Algorithm, and Machine axes. Pruning may reduce parameters while asking the runtime to exploit sparse structure; distillation may reduce inference cost while adding a teacher-student training phase; quantization may reduce data movement while spending numerical precision. The engineer’s task is to move complexity to where the cost is lowest given deployment constraints.

The challenge is removing that surplus without removing what matters. Each technique relocates complexity. Pruning reduces parameter count but asks the hardware to exploit sparse patterns and handle irregular memory access. Knowledge distillation produces a smaller deployment model by spending more compute during training. Neural architecture search reduces human design effort by spending a larger automated search budget. Structural optimization therefore asks where complexity should reside for a given deployment target.³

These three techniques address the challenge through complementary approaches. Pruning eliminates low-impact parameters from an existing model. Knowledge distillation transfers a large model’s learned capabilities to a smaller architecture. NAS automates architecture design from the ground up, building optimized structures for specific constraints (Hutter et al. 2019). In practice, these techniques are often combined: a NAS-designed architecture, distilled from a large teacher, then pruned for final deployment. Pruning comes first because it exposes the central structural trade-off most directly: removing parameters only helps if the resulting structure is something the runtime can exploit.

10.3.1 Pruning

As an illustrative deployment scenario, consider a MobileNet trained for image classification on a wearable health monitor. The trained model occupies 14 MB, but the target microcontroller offers only 2 MB of flash memory. Retraining a smaller architecture from scratch would require weeks of data collection and validation—time the product schedule does not allow. Suppose profiling shows that about 85.7 percent of the model’s weights are near zero and contribute little on the validation set. Removing those weights and fine-tuning the remainder for a few epochs could produce a model that fits in 2 MB with an acceptable accuracy loss. The numbers anchor the engineering trade-off rather than reporting a universal MobileNet benchmark.

Pruning⁴ directly addresses memory efficiency constraints by eliminating parameters or structures that contribute little to the deployed objective. Many trained networks contain removable capacity, but the safe fraction depends on the architecture, task, pruning pattern, and recovery procedure.

² | **Overparameterization:** C. Zhang et al. (2017) demonstrated that networks large enough to fit ImageNet can also memorize completely random labels, showing that training capacity can exceed the structure needed for natural labels. Pruning studies then show the deployment consequence: trained models often contain many parameters that can be removed or sparsified with modest task loss when pruning and fine-tuning are done carefully (Gale et al. 2019; Blalock et al. 2020). The redundancy is not a universal $10\times$ constant; it depends on architecture, dataset, sparsity pattern, and runtime support.

³ | **Pareto frontier:** Named after Italian economist Vilfredo Pareto (1848–1923), who observed that 80 percent of Italy’s land was owned by 20 percent of the population. In multi-objective optimization, the Pareto frontier is the set of solutions where improving one objective (for example, speed) necessarily sacrifices another (for example, accuracy). EfficientNet traces this frontier concretely: B0 (77.1 percent accuracy, 390 million FLOPs) to B7 (84.4 percent, 37 billion FLOPs)—a $95\times$ compute increase for 7.3 percentage points of accuracy, quantifying how steep the trade-off becomes at the frontier’s edge (Tan and Le 2019).

⁴ | **Optimal Brain Damage:** Introduced by LeCun, Denker, et al. (1989), the method achieved $4\times$ parameter reduction—and proportional memory savings—in a handwriting recognizer by using a diagonal approximation to second-derivative (Hessian) information to estimate the loss increase from removing each weight. A full Hessian has $\mathcal{O}(n^2)$ entries for n parameters, but Optimal Brain Damage avoids storing that full matrix by approximating its diagonal. Modern-scale pruning still favors cheaper criteria such as magnitude because even diagonal curvature estimation adds training cost.

The central questions are *what* to prune (individual weights vs. entire structures), *how* to estimate what is expendable (magnitude, gradients, or activations), and *when* to prune (after training, during training, or even at initialization). Section 11.4.5 explains the hardware side; the systems lesson here is that zeros become valuable only when the execution path can skip them.



Definition 10.2: Pruning

Pruning is a model-compression technique that sparsifies the parameter space by removing weights that contribute minimal information to the loss landscape.

1. **Significance:** It can convert dense matrices into sparse structures, reducing memory footprint and total data volume (D_{vol}) when the sparse representation, metadata, and recovery procedure preserve acceptable quality.
2. **Distinction:** Unlike quantization, which reduces the precision of every weight, pruning reduces the count of weights by identifying and eliminating redundancy.
3. **Common pitfall:** A frequent misconception is that pruning “automatically” speeds up execution. In reality, without specialized sparse execution support, the resulting sparse matrices may actually run slower than dense ones due to irregular memory access patterns; a higher R_{peak} alone does not make an irregular sparse layout efficient.

We seek a sparse version of the model parameters $\hat{\mathbf{W}}$ that minimizes the increase in prediction error (loss) while satisfying a fixed parameter budget k . Framing this goal mathematically clarifies both the objective and why approximate solutions are necessary:

$$\min_{\hat{\mathbf{W}}} \mathcal{L}(\hat{\mathbf{W}}) \quad \text{subject to} \quad \|\hat{\mathbf{W}}\|_0 \leq k$$

where $\|\hat{\mathbf{W}}\|_0$ is the **L0-norm** (the count of nonzero parameters). Solving this cardinality-constrained optimization is combinatorial, so practical methods use heuristics⁵ such as **magnitude-based pruning**. Listing 10.1 demonstrates this approach, removing weights with small absolute values to transform a dense weight matrix into the sparse representation visualized in figure 10.2.

Listing 10.1: Magnitude-Based Pruning: Removes weights below a threshold to create sparse matrices, reducing the number of nonzero parameters from 9 to 4 ($k = 4$).

```
import torch

# Original dense weight matrix
weights = torch.tensor(
    [[0.8, 0.1, -0.7], [0.05, -0.9, 0.03], [-0.6, 0.02, 0.4]]
)

# Simple magnitude-based pruning: keep only the 4 largest weights
threshold = 0.5
mask = torch.abs(weights) >= threshold
pruned_weights = weights * mask

print("Original:", weights)
print("Pruned (4 nonzeros):", pruned_weights)
```

Notice how the sparse matrix on the right retains only the high-magnitude values (colored cells) while the near-zero weights become exactly zero. In models amenable to pruning, much of the useful information can remain after many low-magnitude weights are removed. This observation motivates magnitude-based pruning as a practical heuristic.

To make pruning computationally feasible, practical methods often replace the hard L0 constraint with soft regularization like L1-norm ($\lambda_{\text{L1}} \|\mathbf{W}\|_1$), where λ_{L1} controls the strength of the sparsity penalty and $\mathbf{W}$ denotes the weight tensor being regularized. This encourages small values that can

⁵ | **Heuristic:** From Greek *heuriskein* (to discover), the same root as Archimedes’ “eureka.” In pruning, the dominant heuristic—larger magnitude means more important—works well empirically but creates a systems trap: magnitude-based pruning applied globally can remove most parameters from overparameterized layers while leaving critical bottleneck layers largely intact, giving the appearance of aggressive compression while preserving much of the compute and memory cost in the layers that matter (Blalock et al. 2020). This is why iterative prune-retrain cycles with per-layer budgets are often safer than naive global magnitude pruning: each cycle lets the network redistribute importance before the next cut.




Figure 10.2: Sparse Matrix Transformation: Magnitude-based pruning replaces weights below a numerical threshold with exact zeros (white cells) while preserving salient high-magnitude connections (colored cells). The resulting sparsity cuts parameter count, but realizing hardware speedups requires translating these unstructured zero entries into sparse storage formats or specialized execution kernels.

later be thresholded to zero. Practitioners typically use **iterative pruning**, where parameters are removed in successive steps interleaved with fine-tuning to recover lost accuracy (Gale et al. 2019; Blalock et al. 2020).

10.3.1.1 Target structures

The choice of what to prune depends on the deployment target’s hardware constraints and which resource is the binding bottleneck. When memory capacity is primary, **neuron pruning** can provide direct relief in models whose fully connected layers dominate parameter storage: removing entire neurons and their associated weights and biases reduces layer width and parameter count. Profiling must first establish that these layers, rather than embeddings, activations, or another component, are the actual bottleneck.

When convolutional inference latency on commodity accelerators is the bottleneck, **channel pruning** (also called filter pruning) is a strong candidate. Eliminating entire channels or filters reduces feature-map depth and the multiply-accumulate count in subsequent layers. The resulting subnetwork remains dense and regular, so it can map to conventional GPU and Tensor Processing Unit (TPU) kernels without an unstructured sparse format. Realized latency still depends on the resulting dimensions, kernels, and memory traffic.

When a model contains removable stages, **layer pruning** removes entire layers from the network. Each removal eliminates all computation in that stage, but it also reduces representational depth and may change tensor shapes, residual paths, or downstream interfaces. The remaining layers must absorb the lost function, typically through fine-tuning or retraining. The nominal operation saving therefore does not guarantee proportional latency: graph rewrites, kernel shapes, and memory traffic still govern execution. Layer pruning demands careful task validation and end-to-end profiling. The side-by-side comparison in figure 10.3 shows why channel and layer pruning have different implementation costs.

To see how these approaches differ in practice, compare the two sides of figure 10.3. When a channel is pruned, the model’s architecture must be adjusted to accommodate the structural change. Specifically, the number of input channels in subsequent layers must be modified, requiring alterations to the depths of the filters applied to the layer with the removed channel. In contrast, layer pruning removes all channels within a layer, necessitating more significant architectural modifications. In this case, connections between remaining layers must be reconfigured to bypass




Figure 10.3: Channel vs. Layer Pruning: Structural pruning operates at different architectural granularities. Channel pruning (left) removes feature map slices within convolutional layers while preserving overall network topology and dense execution, whereas layer pruning (right) eliminates entire computational stages, requiring graph rewiring and risking greater capacity loss to achieve larger per-operation savings.

the removed layer. Regardless of the pruning approach, fine-tuning is important to adapt the remaining network and restore performance.

10.3.1.2 Unstructured pruning

Unstructured pruning removes individual weights while preserving the overall network architecture. Some connections become redundant during training, contributing little to the final output. Pruning these weak connections reduces the nonzero count and can preserve most task quality after recovery.

Formalizing this process, let $\mathbf{W} \in \mathbb{R}^{m \times n}$ represent a weight matrix in a given layer. Pruning removes a subset of weights by applying a binary mask $\mathbf{M} \in \{0, 1\}^{m \times n}$, yielding a pruned weight matrix:

$$\hat{\mathbf{W}} = \mathbf{M} \odot \mathbf{W}$$

where $\odot$ represents the element-wise Hadamard product. The mask $\mathbf{M}$ is constructed based on a pruning criterion, typically weight magnitude. A common approach is magnitude-based pruning, which removes a fraction ρ_{sparse} of the lowest-magnitude weights by defining a threshold δ_{prune} such that:

$$M_{i,j} = \begin{cases} 1, & \text{if } |W_{i,j}| > \delta_{\text{prune}} \\ 0, & \text{otherwise} \end{cases}$$

where δ_{prune} is chosen to ensure that only the largest $(1 - \rho_{\text{sparse}})$ fraction of weights remain. This method assumes that larger-magnitude weights contribute more to the network's function, making them preferable for retention.

Unstructured pruning reduces the number of nonzero parameters. That reduction saves artifact memory only when the zeros are encoded in a sparse format whose metadata costs less than the omitted values. Otherwise, a dense tensor with masked zeros occupies the original storage.

Unstructured pruning does not necessarily improve computational efficiency on modern hardware, however. Standard accelerators are optimized for dense matrix multiplications, and a sparse weight matrix often cannot fully use hardware acceleration unless specialized sparse computation kernels

are available. Unstructured pruning therefore primarily benefits model storage rather than inference acceleration.

10.3.1.3 Structured pruning

Where unstructured pruning removes individual weights, **structured pruning** (H. Li et al. 2017) eliminates entire computational units: neurons, filters, channels, or layers. This approach can produce smaller dense tensors that use ordinary accelerator kernels. It is therefore often easier to turn into latency savings than arbitrary unstructured sparsity, although the new shapes, memory behavior, and task quality still require measurement.

Neurons, filters, and layers can differ substantially in their contribution to a model's predictions. Some units may carry redundant or low-impact information, but removal is safe only when validation shows that the remaining structure preserves required behavior. Identifying those structures remains the core challenge.

Hardware-aware pruning strategies, such as **N:M structured sparsity**,⁶ enforce specific patterns (for example, ensuring 2 out of every 4 weights are zero) to align with specialized accelerator capabilities. This chapter uses the 2:4 pattern as the compression example; section 11.4.5 later shows how sparse Tensor Cores exploit it.

Pruning methods trade parameter reduction against hardware execution regularity. Compare the three panels in figure 10.4, contrasting irregular weight removal on the left with structured neuron and channel removal in the center and right panels.




Figure 10.4: Unstructured vs. Structured Pruning: The regularity-compression trade-off across pruning granularities. Unstructured pruning (left) removes arbitrary individual connections and requires a sparse representation and matching kernel to change execution. Structured pruning of neurons (middle) or convolutional channels (right) can produce smaller dense tensors that remain compatible with standard kernels, although realized latency depends on the resulting shapes and runtime. Source: (Qi et al. 2021).

In contrast, structured pruning (depicted in the middle and right sections of figure 10.4) removes entire neurons or filters while preserving the network's overall structure. In the middle section, a pruned fully connected network retains its fully connected nature but with fewer neurons. On the right, structured pruning is applied to a CNN by removing convolutional kernels or entire channels (dashed squares). This method maintains the CNN's core convolutional operations while reducing the computational load, making it more compatible with hardware accelerators.

A common approach to structured pruning ranks entire neurons or filters by the magnitude of their associated weights. A low group norm is an inexpensive proxy for low importance, not proof that the unit is dispensable. The score can use an ℓ_1 -norm or ℓ_2 -norm over the weights associated with each unit; groups below a selected threshold become pruning candidates. Layer-wise ranking matters because raw scales can differ across layers. The method needs no representative data pass, but computing scores, choosing a threshold, materializing the smaller graph, and validating or fine-tuning the result still incur engineering and compute cost.

Another strategy is **activation-based pruning**, which evaluates neuron or filter activations over a dataset. Units that remain weak across representative inputs become candidates for removal, subject to ablation or fine-tuning checks. This method captures input-dependent behavior rather than relying only on static weights. Its ranking can still miss rare but important inputs, depend on

⁶ | **N:M structured sparsity:** Introduced commercially with NVIDIA's A100 GPU (2020), the 2:4 pattern was chosen because it halves multiply-accumulate operations while keeping position metadata small enough for the sparse Tensor Core path (NVIDIA Corporation 2020a; Choquette et al. 2021). This fixed ratio is a hardware constraint, not a mathematical optimum: the A100 Sparse Tensor Core path accelerates 2:4 sparse operands, yielding up to 2× math-throughput speedup over dense execution when kernels and layouts satisfy the constraint. Other ratios are not accelerated by this specific hardware path, illustrating how silicon design constrains which sparsity patterns translate to actual speedup.

activation scaling, or change after upstream units are removed. Activation-based pruning therefore requires representative profiling data and validation after the selected structures are removed.

Gradient-based pruning uses loss derivatives to estimate how removing a neuron or filter would change the objective. A small first-order score suggests a candidate at the current checkpoint and on the sampled data; it does not establish permanent irrelevance or capture every interaction among units. Gradient-based criteria can combine a unit's value with its gradient to rank the expected local loss change. Unlike weight-only magnitude criteria, they require backward computation and are usually integrated into training or fine-tuning. The pruned result still needs task-level validation because the score is a local approximation.

These criteria trade measurement cost against fidelity to model behavior. Magnitude scores are inexpensive and provide a useful baseline, but they ignore the input distribution. Activation scores incorporate representative inputs and expose units that are quiet on that workload, while gradient-based scores also incorporate the local loss surface at additional training cost. None dominates across models, and none establishes that removal is harmless. A sound comparison prunes to the same structural target, applies the same recovery budget, and measures task quality after removal. It then profiles the exported graph, because a better importance score can preserve quality without producing shapes that the target runtime executes efficiently. The choice therefore depends on available data and compute, the structure being removed, and the deployment's quality margin.

10.3.1.4 Dynamic pruning

Traditional pruning methods, whether unstructured or structured, produce a **Static pruning** mask or subnetwork that remains fixed during deployment, even if the mask was learned gradually during training. **dynamic pruning** instead adapts which computation is active from input data or training dynamics, so the executed subnetwork can change over time.

Dynamic pruning can use runtime sparsity techniques in which the model selects parameters or structures from input characteristics. Activation-conditioned pruning, for example, deactivates neurons or channels for particular inputs (Hu et al. 2023). The resulting input-dependent sparsity reduces executed work only when the controller and runtime can skip that work more cheaply than a dense path would execute it.

For instance, consider a convolutional neural network processing images with varying complexity. During inference on a simple image containing mostly uniform regions, some convolutional filters may produce negligible activations. A dynamic policy can identify candidate filters and temporarily exclude them from computation. That exclusion reduces nominal work, but the controller, irregular execution, and any quality change remain part of the result. The method becomes useful in a latency-sensitive application only when variation among inputs is large enough to repay those costs. Chapter 12 presents measurement strategies for evaluating such efficiency gains; at this point, the key requirement is to measure both speed and accuracy on the same target workload and across the full input distribution.

Another class of dynamic pruning operates during training, gradually introducing and adjusting sparsity throughout the optimization process. Gradual magnitude pruning starts with a dense network and progressively increases the fraction of pruned parameters while fine-tuning the surviving weights. Dynamic sparse training variants additionally allow the network to recover from pruning-induced capacity loss by regrowing connections that prove important later in training.

Dynamic pruning can allocate different amounts of work to different inputs and can reactivate capacity that a static deployment mask would remove permanently. Those properties create an opportunity, not a guaranteed efficiency or accuracy gain. A controller sits on the critical path, input-dependent routes can fragment batches, and the enlarged behavior space requires more validation. Training may also need routing losses or regularization so that the policy actually uses its cheaper paths. Production deployments must monitor route frequency, latency tails, and quality by route; chapter 14 later develops those monitoring and rollback practices. Dynamic pruning is therefore a candidate when input complexity varies enough to repay the controller and batching overhead, not a default substitute for static pruning.

10.3.1.5 Pruning trade-offs

The three pruning approaches occupy different positions on the regularity-versus-granularity trade-off. Unstructured pruning can remove individual weights and therefore express fine-grained sparsity,

but accelerators need matching sparse formats and kernels to skip those zeros. Structured pruning removes channels, filters, or layers and can produce smaller dense operations that conventional hardware already supports. Dynamic pruning makes the executed structure input-dependent, adding control and batching overhead in exchange for conditional allocation. Table 10.4 summarizes these deployment distinctions.

These categories describe deployed execution rather than mutually exclusive training algorithms. A structured subnetwork may be learned with a dynamic schedule and then frozen for serving; classify the artifact by what its runtime actually executes.

Table 10.4: Pruning Strategies: Unstructured pruning removes individual weights but needs sparse storage and matching kernels for deployment gains; structured pruning removes whole components that map more directly to dense hardware, while dynamic pruning adds runtime adaptation and overhead.

Aspect	Unstructured Pruning	Structured Pruning	Dynamic Pruning
What is removed?	Individual weights in the model	Entire neurons, channels, filters, or layers	Adjusts pruning based on runtime conditions
Model structure	Sparse weight matrices; original architecture remains unchanged	Model architecture is modified; pruned layers are fully removed	Structure adapts dynamically
Impact on memory	Reduces the nonzero count; storage falls only with an appropriate sparse representation	Reduces model storage by removing entire components	Varies based on real-time pruning
Impact on computation	Dense work remains unless a supported sparse path skips zeros	Reduces nominal FLOPs; latency depends on resulting shapes and kernels	Varies with routing policy and controller overhead
Hardware compatibility	Requires sparse formats and matching execution support	Produces dense subnetworks, but awkward shapes can still underutilize hardware	Requires adaptive inference engines
Fine-tuning required?	Often used to recover accuracy after pruning	Often used after structural modification	Depends on how the routing policy and subnetworks are trained
Use cases	Memory-efficient model compression for cloud deployment	Real-time inference optimization, mobile/edge AI, and efficient training	Adaptive AI applications, real-time systems

10.3.1.6 Pruning strategies

Beyond the broad categories of unstructured, structured, and dynamic pruning, different pruning workflows can impact model efficiency and accuracy retention. Two widely used pruning strategies are iterative pruning and one-shot pruning, each with distinct benefits and trade-offs.

Iterative pruning Iterative pruning mitigates drastic accuracy degradation by interleaving structural removal with retraining cycles. Trace the three-row workflow in figure 10.5, following how intermediate accuracy drops recover after each fine-tuning phase.

Follow the three rows of figure 10.5 to see this gradual process in action on a convolutional neural network where six channels are pruned. Rather than removing all channels simultaneously, iterative pruning eliminates two channels per iteration over three cycles. Following each pruning step, the model undergoes fine-tuning to recover performance. The first iteration, which removes two channels, results in an accuracy decrease from 0.995 to 0.971, but subsequent fine-tuning restores accuracy to 0.992. After completing two additional pruning-tuning cycles, the final model achieves 0.991 accuracy, which represents only a 0.4 percent reduction from the original, while operating with 27 percent fewer channels. By distributing structural modifications across multiple iterations, the network maintains its performance capabilities while achieving improved computational efficiency.

One-shot pruning One-shot pruning removes multiple architectural components in a single step, followed by an extensive fine-tuning phase to recover model accuracy. This aggressive approach compresses the model quickly but risks greater accuracy degradation, as the network must adapt to significant structural changes simultaneously.

Consider applying one-shot pruning to the same network from the iterative pruning example. Instead of removing two channels at a time over multiple iterations, one-shot pruning eliminates




Figure 10.5: Iterative Pruning Performance: Three rows depict successive prune-then-fine-tune cycles, each removing two of the original twenty-two channels. Accuracy drops from 0.995 to 0.971 after the first prune, recovers to 0.992 after fine-tuning, and settles at 0.991 after all three cycles, a 0.4 percent loss with 27 percent fewer channels.

all six channels simultaneously. Compare the single-row workflow in figure 10.6 to the iterative case: removing 27 percent of the network’s channels simultaneously causes the accuracy to drop significantly, from 0.995 to 0.914. Even after fine-tuning, the network only recovers to an accuracy of 0.943, which is a 5 percent degradation from the original unpruned network. While both iterative and one-shot pruning ultimately produce identical network structures, the gradual approach of iterative pruning better preserves model performance.




Figure 10.6: One-Shot Pruning Impact: All six channels (27 percent) are removed simultaneously, causing accuracy to drop from 0.995 to 0.914. Fine-tuning recovers only to 0.943, a 5 percent degradation from the original accuracy.

The choice between strategies depends on three interrelated factors. First, the sparsity target: higher reduction targets often necessitate iterative approaches to maintain accuracy, while moderate goals may be achievable with one-shot methods. Second, available resources: iterative pruning demands significant compute for multiple fine-tuning cycles, whereas one-shot approaches trade accuracy for speed. Third, the deployment timeline and target platform: one-shot methods enable

faster deployment, but certain hardware architectures better support specific sparsity patterns, making iterative approaches more advantageous when time permits.

10.3.1.7 Lottery ticket hypothesis

The pruning strategies in this chapter share a common assumption: they start with a trained network and then identify which parameters to remove. The relationship between network structure and trainability may run deeper than pruning strategies suggest: pruning may reveal inherently efficient subnetworks that were already hidden within the dense model, rather than merely deleting unnecessary weights after training.

This perspective leads to the Lottery Ticket Hypothesis⁷ (LTH), which challenges conventional pruning workflows by proposing that within large neural networks, there exist small, well-initialized subnetworks (“winning tickets”) that can achieve comparable accuracy to the full model when trained in isolation. Rather than viewing pruning as a post-training compression step, LTH suggests it can serve as a discovery mechanism to identify these efficient subnetworks early in training (Rachwan et al. 2022).

The original LTH experiments test the hypothesis through an iterative pruning process. Trace the cycle in figure 10.7: a network is trained, low-magnitude weights are pruned, and the survivors are reset to their original initialization rather than re-randomized. Repeating the cycle can isolate a sparse subnetwork that, in the reported settings, trains to accuracy comparable with the original network (Frankle and Carbin 2019). Whether the same procedure succeeds depends on architecture, scale, optimizer, and training budget.




Figure 10.7: Lottery Ticket Iteration Cycle: The original procedure repeatedly trains, prunes low-magnitude weights, and resets surviving weights to their initial state θ_0 . The cycle tests whether the resulting sparse subnetwork can match the dense baseline when trained in isolation; the outcome is empirical rather than guaranteed.

The Lottery Ticket Hypothesis suggests that compact, trainable subnetworks may exist inside an overparameterized initialization. That observation motivates direct sparse-training research, but the standard discovery procedure does not eliminate dense training cost: it first pays for one or more dense or partially dense runs to find the ticket. The result therefore emphasizes the role of initialization while leaving a practical systems question—whether the subnetwork can be identified before most of the original training budget is spent.

The iterative procedure also separates two effects that one-shot pruning conflates: which weights survive and how much recovery training follows each cut. In some settings, smaller pruning steps give the network more opportunity to recover; in others, the additional cycles do not justify their

⁷ | **Lottery ticket hypothesis:** Named for the intuition that training a large network is like buying many lottery tickets—most lose, but a few “winning tickets” (sparse subnetworks with favorable initializations) can train to comparable accuracy on their own. Frankle and Carbin (2019) established the hypothesis on smaller vision and fully connected networks; later work surveyed and extended the idea to larger settings (Rachwan et al. 2022). The systems implication is that some of the memory and compute spent training dense networks may be discoverable overhead, but the practical payoff depends on whether the winning subnetwork can be found before paying most of the original training cost.

cost. The deployment comparison must therefore measure final quality, search or retraining cost, and realized runtime savings together.

In practice, an LTH-style pipeline should be compared with training a smaller dense model from the start, one-shot pruning with recovery, and direct sparse-training methods. The accounting must include every dense or partially dense discovery run, checkpoint rewind, pruning cycle, and recovery phase—not only the final subnetwork. A winning ticket is a training result; it becomes a systems win only when total accelerator time, final artifact size, target latency, and task quality beat those alternatives.

10.3.1.8 Pruning in practice

LTH presents a compelling theoretical perspective on pruning, but practical implementations must still produce runtime artifacts that deployment systems can exploit. Framework pruning is only useful when it produces a deployment artifact the runtime can exploit: a smaller dense model, a structured sparse model, or a sparse format with matching kernels. Training-time pruning often begins as mask application: the original tensor remains present, but a binary mask zeros selected weights during the forward pass. That mechanism can guide learning, yet by itself it does not guarantee lower latency or memory use. Deployment savings appear only after the masked structure is materialized into the artifact that the serving runtime actually loads.

This artifact boundary separates the pruning strategies. Unstructured pruning removes individual weights and needs sparse kernels plus an appropriate storage format to translate zeros into speed. Structured pruning removes whole channels, heads, neurons, or blocks, which can reshape tensors into smaller dense operations that ordinary accelerators already execute well. Gradual pruning during fine-tuning adds a training schedule: sparsity increases over time so the remaining weights can recover accuracy as capacity is removed. The systems audit is therefore concrete: identify what is pruned, identify the runtime format, and verify that the target hardware has kernels that make the sparsity useful.

These trade-offs become concrete when examining real-world deployments. Some model families reduce deployment cost through architecture rather than post-hoc pruning: MobileNet uses depthwise separable convolutions for mobile and embedded vision (Howard et al. 2017), while EfficientNet uses compound scaling to improve the accuracy-efficiency trade-off under resource constraints (Tan and Le 2019). Pruning remains a separate optimization lever. BERT-style transformers⁸ have been pruned by removing redundant attention heads or intermediate dimensions, while separate distillation methods such as DistilBERT and TinyBERT train smaller dense student models that retain much of BERT's performance (Sanh et al. 2019; Jiao et al. 2020).

Pruning has an inherent limitation: it starts with an existing architecture and carves away pieces. The pruned model inherits its structure from the original—same layer types, same connectivity patterns, just fewer parameters. The original architecture itself may be inefficient for deployment. A practitioner may need a model with a completely different structure, such as a six-layer transformer instead of a 12-layer one, that still captures the original model's capabilities.

If deployment requires a different architecture, the next question is how a smaller model can acquire behavior learned by the original. Knowledge distillation answers that question by using the larger model as supervision rather than as the deployment artifact.

10.3.2 Knowledge distillation

Suppose a medical question-answering model exceeds the memory and latency budget of a hospital's single-GPU server. Making its existing weights sparse may still leave an unsupported representation or an architecture too large for the target, so the deployment may require a fundamentally smaller model. Knowledge distillation addresses this class of problem by training a compact "student" to reproduce a larger "teacher's" behavior, often retaining much of the teacher's task performance at lower inference cost (Hinton et al. 2015; Sanh et al. 2019). The term *distillation* borrows from chemistry, where the process extracts a concentrated essence from a larger mixture,⁹ but the systems insight is specific: the teacher's predictions can carry information beyond the raw training labels.

⁸ | **BERT pruning:** BERT's 12 attention heads per layer can contain substantial redundancy—Michel et al. (2019) found on MultiNLI that many heads could be removed without a statistically significant performance change. The removable fraction depends on the layer and task, so deployment pruning must measure both task quality and whether the resulting structure reduces runtime work.

⁹ | **Distillation:** The name borrows from the chemical process of separating a mixture. In model compression, the student learns from a teacher's softened output distribution or internal representations rather than copying its parameters. Soft targets can encode relationships among classes, but the student learns only what its objective and training data expose. The metaphor is useful but not literal: the temperature parameter comes from softmax scaling, and the size-quality trade-off depends on the student, data, objective, and training procedure (Hinton et al. 2015). Deployment savings come from the student architecture, not from the metaphor itself.

Definition 10.3: Knowledge distillation

Knowledge distillation is a model-compression technique that trains a smaller *student* model to match the behavior of a larger, pretrained *teacher* model.

1. **Significance:** Distillation moves cost from repeated inference to an additional training phase. A student can be substantially smaller than its teacher while retaining useful task performance; DistilBERT, for example, reports 40 percent fewer parameters, 60 percent faster inference, and 97 percent of BERT’s language-understanding performance (Sanh et al. 2019). This trade-off is valuable when the student-training cost is amortized across many deployed queries.
2. **Distinction:** Unlike pruning, which removes parameters from an existing architecture, and quantization, which lowers numerical precision, distillation trains a new dense architecture. The student inherits behavior from the teacher’s output distribution or intermediate representations rather than inheriting the teacher’s full parameter count or layer structure.
3. **Common pitfall:** A frequent misconception is that distillation is lossless compression. In reality, the student is bounded by its own capacity, the quality of the teacher, and the match between the distillation data and deployment distribution; a student can faithfully reproduce teacher errors as well as teacher knowledge.

Teacher models transfer semantic knowledge by generating continuous probability distributions across output classes. Examine the probability bar chart in figure 10.8, observing how non-target class probabilities capture inter-class structural similarities.




Figure 10.8: Soft Target Distribution: The teacher’s relative confidence levels indicate which classes are semantically similar (for example, cat vs. dog), providing a much richer supervision signal than a binary “correct” label.

Knowledge distillation evaluates student models against both teacher predictions and ground-truth targets. Trace the dual-path training flow in figure 10.9, noting how soft distillation loss and hard student loss combine to drive parameter optimization.

Operationally, the workflow has four decisions before training begins: choose a teacher with the desired behavior, choose a student whose dense architecture fits the deployment target, run the teacher on calibration or task data to produce soft targets, and select the temperature and loss weight that balance teacher imitation against the hard labels. The validation step then checks more than accuracy. A successful distilled model must preserve calibration, subgroup behavior, and latency on the target hardware, because the student can inherit the teacher’s errors as easily as its useful uncertainty.

10.3.2.1 Distillation mathematics

Starting from the softmax normalization introduced in section 5.3.2.2, we use a **temperature parameter**¹⁰ T_{distill} to soften the probability distribution. The softmax output for class i becomes:

¹⁰ | **Temperature (softmax):** The term comes from the Boltzmann distribution in statistical mechanics. Dividing logits by a temperature above one produces a softer distribution and can expose relative probabilities among non-target classes. At $T_{\text{distill}}=1$, the expression is standard softmax; increasing the temperature reduces logit gaps but does not guarantee that the revealed probabilities help the student. The useful amount of softening depends on logit scale, teacher calibration, student capacity, loss weighting, and the accompanying T_{distill}^2 factor. Temperature is therefore selected empirically rather than fixed to a universal range.




Figure 10.9: Knowledge Distillation Workflow: An input sample passes through both the teacher and the student network. The teacher produces soft labels via temperature-scaled softmax, while the student output is compared against both the soft labels (distillation loss) and the hard labels (student loss).

$$p_i^{(T_{\text{distill}})} = \frac{\exp(z_i/T_{\text{distill}})}{\sum_j \exp(z_j/T_{\text{distill}})}$$

A higher T_{distill} flattens the softmax distribution and can expose **dark knowledge**—relative class relationships such as a teacher assigning more probability to dog than truck for an image of a cat. The useful range is tuned for the model and objective. The total loss $\mathcal{L}_{\text{distill}}$ balances standard cross-entropy with the KL divergence:¹¹

$$\mathcal{L}_{\text{distill}} = (1 - \gamma_{\text{KD}}) \mathcal{L}_{\text{CE}}(\mathbf{p}_{\text{student}}, y) + \gamma_{\text{KD}} T_{\text{distill}}^2 \mathcal{D}_{\text{KL}}(\mathbf{p}_{\text{teacher}}^{(T_{\text{distill}})} \parallel \mathbf{p}_{\text{student}}^{(T_{\text{distill}})})$$

Here $\mathbf{p}_{\text{teacher}}^{(T_{\text{distill}})}$ and $\mathbf{p}_{\text{student}}^{(T_{\text{distill}})}$ are the teacher and student probability distributions computed with T_{distill} , y is the hard label, and $\gamma_{\text{KD}} \in [0, 1]$ weights the hard-label and distillation terms. The factor T_{distill}^2 compensates for the temperature-induced reduction in gradient magnitude used in this formulation. The resulting student must still be evaluated on task quality, calibration, memory, and latency.

Efficiency gains and trade-offs Distillation’s primary deployment advantage over unstructured pruning is that it can produce a smaller *dense* model. A dense student uses conventional kernels rather than requiring a sparse execution path. This does not make every dense student fast: depth, width, sequence length, operators, numerical format, and runtime still determine the executed graph. DistilBERT, for example, reports 97 percent of BERT’s language-understanding performance with 40 percent fewer parameters and 60 percent faster inference.¹² The same principle can pair a compact computer-vision architecture such as MobileNet with teacher supervision (Howard et al. 2017; Hinton et al. 2015). The architecture supplies the execution pattern and distillation supplies the training signal; neither substitutes for measurement. The student may inherit useful teacher behavior, but it can also inherit teacher errors, so the transfer must be validated on the deployment distribution.

Distillation can also be combined with other compression techniques, but the evidence should be tied to the specific technique being used. For pruning, Gordon et al. (2020) show that BERT can be pruned during pretraining and then transferred to downstream tasks, rather than requiring a separate pruning pass for each task. Distilled students can also be paired with pruning or quantization in deployment pipelines, but those combinations need to be validated for the target task and hardware.

The limitations are real, however. Distillation requires teacher inference and student training, which can cost more up front than a brief post-training pruning pass, although an iterative pruning pipeline can also be expensive. Effectiveness depends on teacher quality and student capacity: an undersized student may not reproduce the required behavior, while an oversized student may miss the deployment target. Chapter 12 provides a broader evaluation framework; locally, the decision should compare task quality, model size, latency, and total training cost.

¹¹ | **KL divergence:** Introduced by Solomon Kullback and Richard Leibler in 1951, $\mathcal{D}_{\text{KL}}(p \parallel q)$ quantifies the extra bits needed to encode samples from distribution p using a code optimized for q . The key asymmetric consequence: $\mathcal{D}_{\text{KL}}(\text{teacher} \parallel \text{student})$ penalizes the student heavily for assigning zero probability to teacher-probable outputs, forcing the student to maintain broad coverage of the teacher’s distribution, including low-probability “soft labels” that carry the teacher’s learned uncertainty. Distillation can affect calibration, which must be evaluated separately.

¹² | **DistilBERT:** The reported model has 66 million parameters versus BERT-Base’s 110 million, is 40 percent smaller and 60 percent faster, and retains 97 percent of its language-understanding performance. Absolute memory and latency still depend on numerical format, batch size, sequence length, runtime, and hardware (Sanh et al. 2019).

Table 10.5 contrasts the key trade-offs between knowledge distillation and pruning across accuracy retention, training cost, inference speed, hardware compatibility, and implementation complexity. DistilBERT and MobileBERT demonstrate architecture redesign plus distillation; pruning can be combined with distillation in other optimization pipelines, but these models should be understood primarily as dense student-model examples.

Table 10.5: Distillation vs. Pruning: Distillation spends additional training and student-design effort to produce a smaller dense model; pruning modifies an existing model, but sparse results need compatible storage and kernels to realize deployment gains.

Criterion	Knowledge Distillation	Pruning
Accuracy retention	Model- and student-dependent; often strong with a capable teacher	Model- and sparsity-dependent
Training cost	Higher – Requires teacher inference and student training	Often lower – Requires pruning and usually fine-tuning
Inference speed	Often favorable for a smaller dense student	Depends – Structured pruning is efficient, unstructured needs special support
Hardware compatibility	High – Works on standard accelerators	Limited – Sparse models may need specialized execution
Implementation effort	More involved: teacher inference, student objective, and training	Varies: choose a criterion, prune, fine-tune, and export

Knowledge distillation is frequently used alongside pruning and quantization for deployment-ready models. How distillation interacts with these complementary techniques determines the effectiveness of multi-stage optimization pipelines.

Pruning modifies an existing parameterization, while distillation transfers behavior into a chosen student architecture. Neither directly represents a dense weight tensor through low-rank factors. For an illustrative 4096×4096 matrix, a rank-128 factorization stores 6.25 percent as many scalar parameters as the dense matrix; whether that approximation is accurate depends on the matrix’s singular-value spectrum. Structured approximation methods exploit this form of mathematical redundancy.

10.3.3 Structured approximations

Structured approximation serves a different deployment decision from pruning or distillation: when a layer’s weights are mathematically redundant, it may be cheaper to represent that redundancy directly than to store the original dense tensor. These methods decompose large weight matrices and tensors into lower-dimensional components because high-dimensional representations often admit compact, low-rank approximations. Low-rank factorization and tensor decomposition offer complementary strategies for achieving this compression.

10.3.3.1 Low-rank factorization

Low-rank matrix factorization (LRMF) approximates weight matrices with lower-rank representations. Given a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$, LRMF finds matrices $\mathbf{U} \in \mathbb{R}^{m \times k}$ and $\mathbf{V} \in \mathbb{R}^{n \times k}$ such that:

$$\mathbf{A} \approx \mathbf{U}\mathbf{V}^T$$

where $k \ll m, n$ is the approximation rank. This is typically computed via singular value decomposition (SVD),¹³ retaining only the top k singular values.

The deployment question is whether the reduction in stored weights and matrix-vector work is large enough to justify a second factor operation.

Systems Perspective 10.2: The bandwidth-compute trade-off

Low-rank factorization can reduce both arithmetic and bandwidth demand when the retained rank is small, though it adds a second factor operation and possible kernel-launch overhead.

¹³ | **Singular value decomposition (SVD):** The Eckart-Young theorem (1936) proves that retaining the top k singular values yields an optimal rank- k approximation in the Frobenius and spectral norms. The key systems trade-off is whether the high, one-time compute cost of this factorization— $\mathcal{O}(mn \cdot \min(m, n))$ —is amortized by the memory and bandwidth savings from using the smaller model in repeated inference calls.

Storing a 4096 by 4096 matrix requires 67.1 MB (at FP32). Fetching this matrix for a single inference imposes substantial memory traffic, especially when execution is limited by physical memory bandwidth.

If we factorize it with rank $k = 128$, we store two matrices (4096 by 128 and 128 by 4096), totaling only 4.2 MB, a 16 $\times$ reduction in stored values. When inference uses the two factors directly, matrix-vector arithmetic also falls from $\mathcal{O}(mn)$ to $\mathcal{O}(k(m+n))$ for small k . The deployment speeds up only if the two factor operations and their intermediate traffic cost less than the original dense kernel. Explicitly materializing UV^T would add $\mathcal{O}(mkn)$ work and defeat the purpose.

This bandwidth-compute trade-off is the local version of the memory wall: execution becomes bottlenecked by moving weights rather than multiplying them. Section 11.5.1 later examines the same phenomenon from the hardware side.

Large weight projection matrices can be compressed into low-rank representations without full rank materialization. Locate the matrix decomposition in figure 10.10, observing how full matrix $\mathbf{M}$ splits into rank- k factors $\mathbf{L}_k$ and $\mathbf{R}_k^T$.




Figure 10.10: Low-Rank Factorization: A weight matrix $\mathbf{M}$ of size $m \times n$ is approximated by two smaller factors, labeled $\mathbf{L}_k$ ($m \times k$) and $\mathbf{R}_k^T$ ($k \times n$) in the diagram and corresponding to $\mathbf{U}_k$ and $\mathbf{V}_k^T$ in the surrounding notation. Storage falls from $m \times n$ to $m \times k + k \times n$ parameters, while one dense operation becomes two factor operations.

LRMF applies directly to fully connected layers and can be applied to suitably reshaped convolutional weight tensors. Storage reduces from $\mathcal{O}(mn)$ to $\mathcal{O}(mk + kn)$, while inference executes two factor operations. Choosing rank k balances compression and arithmetic savings against approximation error and kernel overhead.

10.3.3.2 Tensor decomposition

Tensor decomposition extends factorization to multi-dimensional tensors common in convolutional layers and attention mechanisms, so the deployment decision becomes whether the storage saved by a low-rank tensor representation outweighs the reconstruction and inference overhead. Figure 10.11 breaks down a 3D tensor into its factor matrices, showing how each rank-one component contributes to the reconstruction. The choice among decomposition methods depends on tensor order, target rank, and inference overhead.

The main tensor-decomposition families differ in representation. **CP decomposition** expresses a tensor as a sum of rank-one components, $\mathcal{X} \approx \sum_{r=1}^k \mathbf{u}_r \otimes \mathbf{v}_r \otimes \mathbf{w}_r$ (Lebedev et al. 2015). **Tucker decomposition** keeps a small core tensor with factor matrices, $\mathcal{X} \approx \mathcal{G} \times_1 \mathbf{U} \times_2 \mathbf{V} \times_3 \mathbf{W}$. **Tensor-train (TT)** represents a high-order tensor as a sequence of lower-order cores. Each method's usefulness depends on rank choice, parameter savings, approximation error, and available kernels.

Tensor decomposition applies to convolutional filters (approximating 4D weight tensors), attention mechanisms in transformers, and embedding layers in natural language processing (NLP) models. The trade-offs mirror LRMF: compression vs. information loss, and the additional computational overhead of tensor contractions during inference. Table 10.6 compares LRMF and tensor decomposition across applicable data structure, compression mechanism, and computational cost.




Figure 10.11: Tensor Decomposition: A three-dimensional tensor $y \in \mathbb{R}^{M \times N \times T}$ is approximated by factor matrices U , X , and V . The highlighted factor vectors u_i , x_t , and v_j combine across rank R to reconstruct the entry y_{ijt} . Extending low-rank approximation from two dimensions to three carries the storage and computation argument into convolutional layers. Source: (Gholami et al. 2022).

Table 10.6: Matrix and Tensor Factorization: LRMF applies to two-dimensional matrices, while tensor decompositions extend low-rank representations to multidimensional tensors; realized storage and latency depend on rank choice and kernel support.

Feature	Low-Rank Matrix Factorization (LRMF)	Tensor Decomposition
Applicable Data Structure	Two-dimensional matrices	Multi-dimensional tensors
Compression Mechanism	Factorizes a matrix into two or more lower-rank matrices	Decomposes a tensor into multiple lower-rank components
Common Methods	Singular Value Decomposition (SVD), Alternating Least Squares (ALS)	CP Decomposition, Tucker Decomposition, Tensor-Train (TT)
Compute cost	Depends on the factorization method, selected rank, and kernels	Depends on the decomposition, selected ranks, optimization procedure, and tensor-contraction kernels
Storage Reduction	Reduces storage from $\mathcal{O}(mn)$ to $\mathcal{O}(mk + kn)$	Depends on the decomposition and selected ranks; stores factor matrices and, for some methods, a core tensor
Inference execution	Replaces one dense operation with two factor operations	Introduces tensor contractions whose latency depends on available kernels
Primary Use Cases	Fully connected layers, embeddings, recommendation systems	Convolutional filters, attention mechanisms, multi-modal learning
Implementation Complexity	Easier to implement, often involves direct factorization methods	More complex, requiring iterative optimization and rank selection

In practice, LRMF and tensor decomposition can be combined: fully connected layers compressed via LRMF while convolutional kernels use tensor decomposition. The choice depends on the model’s structure and whether memory or latency is the primary constraint.

Pruning and factorization modify existing model representations, while distillation transfers behavior into a chosen student. Neural architecture search takes a different approach: discovering architectures that are efficient by construction.

10.3.4 Neural architecture search

Pruning, distillation, and factorization begin from architectural choices made by people. Selecting a competitive configuration can require extensive experimentation, and a manually explored design space may omit useful candidates (Elsken et al. 2019). **Neural architecture search (NAS)** automates part of this process by exploring a specified space of architectures and objectives. Early reinforcement-learning NAS optimized validation accuracy (Zoph and Le 2016), weight-sharing methods reduced candidate-evaluation cost (Pham et al. 2018), and hardware-aware NAS added device latency or platform efficiency to the objective (Tan et al. 2019).

14 | Hardware-aware NAS: Optimizes measured latency rather than FLOPs, which can diverge by 3–5× when memory access or operator support dominates. Tan et al. (2019) fed device latency into search and found architectures 1.8× faster than MobileNetV2 at higher accuracy.

The three-stage feedback loop in figure 10.12 shows how NAS works. NAS¹⁴ defines a search space of architectural components and constraints, applies a strategy such as reinforcement learning (Zoph and Le 2016), evolutionary search, or gradient-based optimization, and evaluates candidates against accuracy and efficiency objectives. Each evaluation guides the search toward more promising regions of the architecture space. This process can discover competitive architectures while reducing the amount of manual exploration required.




Figure 10.12: Neural Architecture Search Flow: An automated optimization loop for architecture discovery. The search strategy samples candidate subnetworks from a predefined search space, while the performance estimation engine evaluates candidates against target accuracy and hardware constraints (such as latency or energy), feeding reward signals back to steer subsequent exploration toward Pareto-optimal designs.

10.3.4.1 The NAS optimization problem

The effectiveness of NAS depends on three design decisions: what architectures to search over (the search space), how to explore that space efficiently (the search strategy), and how to evaluate each candidate’s fitness for deployment. The optimization problem begins with a chicken-and-egg constraint: we cannot know how good an architecture is until we train it, but training is expensive. This creates two nested decisions: choosing which operations to include (the architecture) and finding the best parameters for those operations (the weights). The architecture defines *what* to optimize; the weights define *how well* that architecture can perform.

NAS is therefore a **bi-level optimization problem**:¹⁵ the outer loop searches the architecture space $\mathcal{A}$, while the inner loop trains candidate architectures to evaluate performance. Formally, we seek the optimal architecture α^* that minimizes validation loss $\mathcal{L}_{\text{val}}$ under constraints C (latency, memory):

$$\alpha^* = \arg \min_{\alpha \in \mathcal{A}} \mathcal{L}_{\text{val}}(\theta^*(\alpha), \alpha) \quad \text{subject to} \quad C(\alpha) \leq C_{\text{max}}$$

where $\theta^*(\alpha)$ represents the optimal model parameters for architecture α , obtained by minimizing training loss:

$$\theta^*(\alpha) = \arg \min_{\theta} \mathcal{L}_{\text{train}}(\theta, \alpha)$$

The core challenge is the cost of the inner loop: evaluating each candidate requires expensive training. A search space with just 10 choices across 20 layers yields 10^{20} architectures, making exhaustive search impossible. Efficient NAS methods address this by restricting the search space, using faster search strategies, or accelerating evaluation.

10.3.4.2 Search space design

The search space defines what architectures NAS can discover. Well-designed search spaces incorporate domain knowledge to focus search on promising regions while remaining flexible enough to discover novel patterns.

Rather than searching entire network architectures, many NAS systems search for reusable computational blocks, or cells, that can be stacked to form complete networks. A convolutional cell might choose from operations such as 3×3 convolution, 5×5 convolution, depthwise separable convolution, max pooling, or identity connections. A simplified cell with four nodes and two operations per edge yields roughly 10,000 possible cell designs, far more tractable than searching full architectures. NASNet exemplifies this approach, discovering reusable normal and reduction cells that can be stacked to form complete networks across different model sizes.

15 | Bi-level optimization: A formulation where one optimization problem sits inside another. In NAS, the outer level selects an architecture while the inner level trains that candidate’s weights, so early methods paid a full training cost for every architecture evaluated. This nesting is why early NAS required 22,400 GPU-days; weight-sharing methods amortize one training run across many candidates, reducing search cost by roughly 1,000×.

The same search-space decision can incorporate deployment constraints directly. Hardware-aware NAS treats latency, memory, or energy on the target platform as first-class objectives rather than optimizing only for accuracy and FLOPs (Zhang et al. 2020). MobileNetV3 was tuned to mobile-phone CPUs through hardware-aware NAS complemented by NetAdapt (Howard et al. 2019). This hardware-aware approach targets measured deployment behavior rather than FLOP count alone.

10.3.4.3 Search strategies

Search strategies determine how to explore the architecture space efficiently without exhaustive enumeration. Table 10.7 compares the trade-offs between search cost, architectural diversity, and optimality guarantees for each approach.

Table 10.7: NAS Search Strategy Comparison: Representative trade-offs between search efficiency, use cases, and limitations for different NAS approaches. The costs are method-specific results from early reinforcement-learning and evolutionary searches and from DARTS, not universal ranges for each strategy (Zoph and Le 2016; Real et al. 2019; Liu et al. 2019).

Strategy	Search Efficiency	When to Use	Key Challenge
Reinforcement Learning	22,400 GPU-days	Novel domains, unconstrained search	High computational cost
Evolutionary Algorithms	3,150 GPU-days	Parallel infrastructure available	Requires large populations
Gradient-Based (DARTS)	1–4 GPU-days	Limited compute budget	May converge to suboptimal local minima

Reinforcement learning-based NAS treats architecture search as a sequential decision: a controller generates architectures and receives an accuracy reward. The controller (typically a long short-term memory network) learns to propose better architectures through policy gradient optimization. This approach discovered high-performing architectures like NASNet, but its inner loop is expensive because every reward requires training a candidate architecture; Zoph and Le (2016) evaluated roughly 12,800 architectures, totaling 22,400 GPU-days. This cost pushed practical NAS toward weight sharing and lower-cost performance predictors.

Evolutionary algorithms maintain a population of candidate architectures and iteratively apply mutations (changing operations, adding connections) and crossover (combining parent architectures) to generate offspring. Fitness-based selection retains high-performing architectures for the next generation, so useful components such as skip connections or depthwise separable convolutions can be recombined rather than rediscovered from scratch. AmoebaNet used evolution to achieve state-of-the-art results after 3,150 GPU-days (Real et al. 2019), showing both the value of population search and the continuing need for proxy tasks, weight sharing, or massive parallelism to control search cost.

Gradient-based methods like DARTS (Differentiable Architecture Search) (Liu et al. 2019) represent the search space as a continuous relaxation where all possible operations are weighted combinations. Rather than discrete sampling, DARTS optimizes architecture weights and model weights jointly using gradient descent. By making the search differentiable, DARTS reduces search cost from hundreds to just one to four GPU-days, though the continuous relaxation may miss discrete architectural patterns that discrete search methods discover.

Hardware-aware NAS moves beyond FLOPs as a proxy for efficiency, directly optimizing for actual deployment metrics. MnasNet’s search incorporates a latency prediction model trained on thousands of architecture-latency pairs measured on actual mobile phones. The search objective combines accuracy and latency through a weighted product:

$$\text{Reward}(\alpha) = \text{Accuracy}(\alpha) \times \left(\frac{L_{\text{lat,target}}}{L_{\text{lat}}(\alpha)} \right)^\beta$$

where $L_{\text{lat}}(\alpha)$ is measured latency, $L_{\text{lat,target}}$ is the latency constraint, and β controls the accuracy-latency trade-off. This formulation penalizes architectures that exceed latency targets while rewarding those that achieve high accuracy within the budget. In the reported MnasNet search space, varying inverted-residual expansion ratios produced a stronger accuracy-latency trade-off than uniform expansion, illustrating how search can expose nonuniform designs that a narrower manual sweep might not test.

10.3.4.4 When to use NAS

Neural architecture search can discover architectures that outperform hand-designed alternatives, but its computational cost demands careful consideration of when the investment is justified. NAS can be worthwhile for novel hardware platforms with unusual constraints, such as new accelerator architectures or highly constrained edge devices, where existing architectures are poorly optimized. It can also make sense at deployment scales where small efficiency improvements justify the upfront search cost, or when architecture families for cloud, edge, and mobile deployments can amortize one search across many variants.

Conversely, custom NAS is difficult to justify when standard deployment constraints already have well-optimized architecture families. If the compute budget is only a few GPU-days, large reinforcement-learning or evolutionary searches are impractical; differentiable or weight-sharing methods such as DARTS may remain feasible, but still require deployment-scale and validation-cost justification. Rapidly changing requirements also weaken the case because the target can move before search and validation finish.

For most practitioners, existing NAS-discovered or NAS-assisted architectures such as EfficientNet (Tan and Le 2019), MobileNetV3 (Howard et al. 2019), and MnasNet (Tan et al. 2019) provide strong baselines before funding a search from scratch. Their transfer quality still depends on task and hardware. Reserve custom NAS for unusual constraints or deployment scales that can amortize the search and validation investment.

10.3.4.5 Architecture examples

NAS studies have surfaced reusable design patterns within their search spaces. EfficientNet jointly scales depth, width, and resolution with compound coefficients and reports improved accuracy-efficiency trade-offs across its model family (Tan and Le 2019). MobileNetV3 combines hardware-aware search and NetAdapt for phone latency targets (Howard et al. 2019), while FBNet incorporates device-specific latency into a mobile-CPU search objective (B. Wu et al. 2019). These results show what search found under particular spaces, objectives, and measurements; they do not establish that automated search universally outperforms careful manual design.

Beyond convolutional networks, NAS has been applied to transformer architectures, including searches for compact language and vision models under compute or memory constraints. The common systems lesson is narrower than a claim that search always beats manual design: when latency, memory, or energy enters the objective and is measured on the target platform, the search evaluates candidates against the resource that deployment actually constrains. The result still depends on the search space, proxy quality, optimization budget, and final retraining.

The structural techniques covered so far (pruning, distillation, factorization, and NAS) all optimize *what* computations the model performs, including which parameters exist, which connections remain, and how the architecture is structured. These techniques can substantially reduce parameter counts and theoretical FLOPs. Regardless of structural efficiency, every surviving weight and activation must still be stored and processed at some numerical precision.



Checkpoint 10.2: Choosing a structural method

Structural methods relocate deployment cost in different ways.

- ☐ **Structured vs. unstructured pruning:** compare how each pattern interacts with GPU and TPU execution.
- ☐ **Distillation vs. pruning:** explain why a dense student may preserve accuracy better than aggressive pruning while requiring additional training.
- ☐ **NAS selection:** identify when hardware-specific constraints and deployment scale justify architecture search rather than manual design.

This brings us to the second dimension of our optimization framework: the precision at which surviving computations are performed. Structural methods can reduce aggregate storage and operation count; precision optimization changes the bytes and arithmetic associated with each represented value. A 32-bit floating-point value occupies 4 bytes, whereas an 8-bit integer occupies

1 byte before scale or packing metadata. For bandwidth-bound inference, the smaller representation can reduce weight traffic and latency when the runtime maps the model to efficient low-precision kernels. The accuracy cost can be small in some well-calibrated INT8 workloads, but it must be measured for the target model, task, and hardware (Jacob et al. 2018; Gholami et al. 2022).

Quantization is often an attractive first deployment experiment because it leaves the architecture intact and can be applied post-training. Its leverage is greatest when weight or activation traffic binds and the target provides efficient low-precision kernels; otherwise, it may shrink the artifact without accelerating the request.

10.4 Quantization and Precision

The framework section established the gap that compression must close: a 7-billion parameter language model in FP16 needs 14 GB, while the smartphone target offers only 8 GB of shared RAM. Pruning can exceed 70 percent sparsity, but zeros do not shrink a densely stored deployment artifact (Han, Mao, et al. 2016). A complementary lever is reducing the number of bits used to represent each parameter, and the open question this section answers is how far precision can be cut before accuracy collapses. Quantization, the process of reducing numerical precision, offers one of the most impactful optimizations for deployment, because it trades bits for speed and efficiency with model-dependent accuracy loss.

Definition 10.4: Quantization

Quantization is a model-compression technique that reduces information fidelity by mapping high-precision continuous values to a lower-precision discrete set.

1. **Significance:** FP32-to-INT8 conversion reduces the raw payload per quantized value by $4\times$. Real artifact size, traffic, and latency also depend on scales, metadata, packing, operator coverage, and the target kernels.
2. **Distinction:** Unlike pruning, which reduces the count of parameters, quantization reduces the bit depth of selected weights and/or activations.
3. **Common pitfall:** A frequent misconception is that quantization is just “rounding.” In reality, it is a lossy mapping that requires careful range estimation and, often, quantization-aware training (QAT) to minimize its impact on accuracy.

Quantization¹⁶ concerns every neural network weight and activation stored at some numerical precision: FP32 (32 bits), FP16 (16 bits), INT8 (8 bits), or lower. Bit width therefore becomes a critical systems parameter. It constrains raw payload size and cache footprint, affects bandwidth demand when values remain packed, and shapes the area and throughput trade-offs of specialized matrix engines (section 11.4.8).

Three system properties can change. The raw weight payload falls by $4\times$ from FP32 to INT8, before scale and packing metadata. Weight traffic can fall by a similar factor when values remain packed through the memory path, accelerating bandwidth-bound inference such as some low-batch large language model (LLM) decoding. Compute cost can also fall when the target provides efficient INT8 kernels; unsupported operators, conversions, or fallback paths can erase part of the gain (Gupta et al. 2015; Y. E. Wang et al. 2019).

The accuracy cost varies by model, task, quantizer, calibration data, and bit width. Many CNN inference models tolerate INT8 well; transformers, speech models, and outlier-heavy layers may need different granularity, mixed precision, or training-aware methods. The main approaches are **post-training quantization** (PTQ), **quantization-aware training** (QAT), and **extreme quantization** for lower-bit deployment. They differ in training cost and control, not in a guaranteed ordering of accuracy.

The viability of each approach depends on how much precision a particular model can shed before quality deteriorates. Figure 10.13 illustrates a common qualitative pattern: a plateau where modest precision reduction has little measured effect, followed by a model-dependent cliff. The bit width at either boundary is not universal.

¹⁶ | **Quantization:** Rooted in Shannon’s theory of representing continuous signals with discrete values (Shannon 1948), reducing FP32 to INT8 collapses over four billion representable values to just 256. Neural networks often tolerate this because trained weights concentrate information in relative magnitudes, not absolute precision: INT8 inference can stay close to the full-precision baseline with appropriate calibration or quantization-aware training, while INT4 and lower-bit methods become more architecture- and method-dependent (Jacob et al. 2018; Gholami et al. 2022; Shen et al. 2020; Lin et al. 2024). The systems consequence is that quantization viability must be validated per-model and per-task, not assumed from aggregate benchmarks.




Figure 10.13: Illustrative Quantization Regimes: The curves depict a qualitative plateau followed by an accuracy cliff as bit width falls. Their breakpoints are illustrative; actual boundaries depend on the model, layer distributions, quantizer, calibration or training method, and evaluation task.

10.4.1 Precision and energy

Precision is an energy decision as much as a storage decision. Efficient numerical representations reduce storage requirements, computation latency, and power usage, benefiting mobile AI, embedded systems, and cloud inference alike. Precision levels can be tuned to specific hardware capabilities, maximizing throughput on AI accelerators such as GPUs, TPUs, NPUs, and edge AI chips.

🔍 Systems Perspective 10.3: The physics of quantization

The energy-movement invariant ($E_{\text{move}} \gg E_{\text{compute}}$) means that, in the physics of silicon, every fetched bit carries an energy cost (Horowitz 2014). Section D.1 collects the reference energy ratios, and section D.4 compares numerical format trade-offs.

According to the iron law ($T = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}}$), reducing bit width changes the data-movement and compute terms only through the execution path that uses the new representation. A first-order packed-value model gives two useful comparisons:

1. **Memory movement** ($D_{\text{vol}} \times E_{\text{move}}$): A packed INT8 stream carries one quarter as many raw value bits as FP32. The corresponding energy estimate assumes full transactions are amortized across packed values; caches, metadata, alignment, and transaction size affect the realized cost.
2. **Compute work** ($O \times E_{\text{compute}}$): A 32-bit multiply costs ≈ 3.7 pJ/op. An INT8 multiply costs ≈ 0.2 pJ/op.

Table 10.8 normalizes these costs against an 8-bit integer add to expose the four-order-of-magnitude gap between arithmetic and DRAM access.

For inference workloads, moving from FP32 to INT8 saves $4\times$ in raw value storage, while the two multiply costs above give an idealized multiply-energy reduction of $18.5\times$. End-to-end energy also includes memory, control, sensors, and idle power. Section D.2.1 gives this relationship a formal shape: workloads with little arithmetic per byte are memory-bound, so shrinking data movement dominates; workloads with more arithmetic per byte are compute-bound, so arithmetic savings matter more. The arithmetic ratio alone cannot predict device battery life.

These same physics apply at data center scale: distributed training systems use reduced precision to cut gradient communication overhead, a topic covered in section 8.6.3. Chapter 11 returns to the silicon mechanisms that exploit these energy differences.

To understand *why* numerics matter so deeply, move from the algorithm to silicon-level energy and data movement. At that level, each bit is both a storage choice and an energy cost.

These savings explain why systems exploit quantization, not why neural networks tolerate it. Numerical tolerance depends on model- and layer-specific robustness to perturbation, while the hardware benefit still depends on preserving task accuracy. How much precision a model can shed before accuracy collapses, and how large the resulting speedup proves to be, are the questions the quantization section develops in full.

10.4.1.1 Energy costs

Table 10.8 summarizes the canonical gap between arithmetic and data movement; precision granularity extends it. Figure 10.14 gives representative operation-level energy costs across more formats and across the SRAM hierarchy: a 32-bit integer addition costs 0.1 pJ/op, while an 8-bit integer addition is just 0.03 pJ/op. Floating-point arithmetic follows the same direction: a 32-bit floating-point addition consumes approximately 0.9 pJ/op, whereas a 16-bit floating-point addition requires 0.4 pJ/op. Relative to the 0.03 pJ INT8 add, the 5 to 50 pJ SRAM reads shown cost approximately 167 to 1,667 times as much energy. This tenfold spread across SRAM sizes shows why where an operand resides can matter as much as how it is computed. DRAM access is more expensive still. These savings compound across large-scale models operating over billions of operations. Reducing a model's energy footprint therefore contributes to sustainable and accessible AI deployment: it helps mitigate the environmental impact of large-scale ML training and inference as AI workloads scale (Patterson et al. 2021), and it expands the reach of machine learning into low-resource environments, from rural healthcare to autonomous systems operating in the field.

Table 10.8: The Energy Hierarchy: Relative energy cost of representative arithmetic and memory operations, normalized to an 8-bit integer add. A 32-bit float add costs 30 $\times$ more than its 8-bit integer counterpart, but a single 32-bit DRAM read costs 21,333.3 $\times$ more, placing DRAM access more than four orders of magnitude above on-chip compute. This gap is why bit-width reductions that shrink data movement dominate the energy budget of inference far more than the arithmetic savings alone would suggest (Horowitz 2014).

Operation	Bit-Width	Relative Energy
Integer Add	8-bit	1 $\times$
Float Add	32-bit	30 $\times$
DRAM Read	32-bit	21,333.3 $\times$


10.4.2 The energy dividend: INT8 vs. FP32

Memory reduction is often the first motivation for quantization, but lower-precision arithmetic can also provide an **energy dividend**. Moving from FP32 to INT8 reduces raw value payload by exactly 4 $\times$; the operation-level energy ratio can be larger, while end-to-end energy remains workload- and device-dependent.

The framework section's energy hierarchy already exposed this asymmetry. In the reference technology model used here, an FP32 addition costs 0.9 pJ/op, while an INT8 addition costs 0.03 pJ/op, an operation-level ratio of 30 $\times$. That ratio is not a device-level battery-life prediction: memory accesses, accumulation precision, control, sensors, and idle power remain. It does explain why many accelerators add specialized low-precision units and why quantization can be necessary when arithmetic or memory energy binds.

These energy savings take on a different character for models where memory capacity, not compute, is the binding constraint.

Recommendation models turn the same energy argument into a placement problem: quantization must make the dominant memory object small enough for the machine that serves it.



Quantization pays off only when the machine axis has the right integer units.




Figure 10.14: Energy per Operation by Precision: Representative 45 nm energy costs range from 0.03 pJ for an INT8 add to 3.10 pJ for an INT32 multiply and 5 to 50 pJ for the SRAM reads shown. The diagonal marks the 8 KB read at approximately 167 times the INT8-add energy; the complete SRAM range spans approximately 167 to 1,667 times that baseline. Energy figures from (Horowitz 2014).

 Lighthouse 10.1: DLRM and embedding quantization

The DLRM lighthouse (section 6.7) presents a compression challenge where memory capacity, rather than compute throughput or memory bandwidth, is the binding constraint (Naumov et al. 2019). Its embedding tables can reach terabytes in size, far exceeding GPU memory.

For DLRM, quantization is not about faster math; it is about storage density. Reducing embedding-table precision from FP32 to INT8 (or lower) can reduce memory footprint by 4–8×, allowing larger tables to fit on fewer GPUs when accuracy and lookup kernels tolerate the lower precision. This is a pure information-density optimization: we compress the lookup table so the *machine* (physics) can hold the *algorithm* (logic).

Together, DLRM and the following TinyML lighthouse separate the two reasons quantization matters. DLRM uses lower precision to make enormous tables fit; TinyML uses lower precision to keep every inference within a tiny energy and SRAM budget. Beyond direct compute savings, reducing numerical precision also lowers memory energy consumption, which often dominates total system power. Lower-precision representations reduce data storage requirements and memory bandwidth usage, leading to fewer and more efficient memory accesses. Accessing memory, particularly off-chip DRAM, is far more energy-intensive than performing arithmetic operations: the representative 32-bit DRAM read used in this chapter costs 640 pJ, compared with picojoule-scale cache and arithmetic operations. An instruction’s total energy can therefore be dominated by memory access patterns rather than computation.¹⁷

 Lighthouse 10.2: The TinyML quantization imperative

The keyword spotting (KWS) lighthouse, DS-CNN, operates at the opposite extreme of the iron law from DLRM (Y. Zhang et al. 2017). A typical TinyML budget is roughly 512 KB of SRAM, and even the purpose-built DS-CNN keyword spotter overshoots it in FP32: its weights occupy 800 KB. Quantized to INT8, the same network shrinks to 200 KB and fits with room for runtime buffers. A stricter smart-doorbell microcontroller with 256 KB SRAM leaves closer to 100 KB for model weights after runtime buffers, audio windows, and feature extraction state, motivating more aggressive INT4 or binary compression.

17 | INT8 energy impact: The energy dominance of memory access is extreme: with the Horowitz constants used here, a single 32-bit DRAM read costs roughly 2,782.6× the energy of an INT8 multiply-accumulate (Horowitz 2014). Quantizing from FP32 to INT8 attacks this disparity on both fronts—4× fewer bytes moved and cheaper arithmetic per operation—although realized energy savings depend on the memory hierarchy, kernels, and workload.

In FP32, even the compact DS-CNN architecture moves $4\times$ more raw weight bytes than in INT8, and the arithmetic path can also cost more on hardware with efficient integer units. For an always-on coin-cell device, that reduction contributes to longer battery life, but sensor, wake-up, feature-extraction, and idle-power costs may dominate. Here, quantization reduces the model's data-movement term and can improve its compute term ($O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$) when the integer path is supported.

Reducing numerical precision thus improves efficiency on two fronts: faster computation and less data movement. This dual benefit is especially valuable for hardware accelerators and edge devices, where memory bandwidth and power efficiency are binding constraints.

10.4.2.1 Performance gains

Figure 10.15 compares illustrative FP32 and INT8 alternatives. The paired bars are not common-platform latency measurements; the storage panel shows the exact $4\times$ reduction from 32 to 8 bits.




Figure 10.15: Illustrative Quantization Scenario: Paired bars separate FP32 and INT8 alternatives. The right panel shows raw weight storage, where reducing each weight from 32 to 8 bits produces an exact fourfold decrease; the left panel shows illustrative, architecture-dependent latency changes.

To make these gains concrete, consider the quantization savings when deploying a modern large language model at reduced precision.

Napkin Math 10.1: Quantization savings

Problem: Deploying Llama 3 8B requires storing 8B parameters on-device. How much memory does the model consume at FP16 vs. INT4, and does quantization shrink it enough to fit on a single consumer GPU?

FP16

- **Raw weight size:** $8 \times 10^9 \times 2$ bytes (16-bit) = 16 GB
- **Runtime budget:** 16 GB leaves no headroom; a 24 GB-class device is practical.

INT4

- **Raw weight size:** $8 \times 10^9 \times 0.5$ bytes (INT4) = 4 GB
- **Runtime budget:** 4 GB leaves about 4 GB on an 8 GB device for metadata, cache, activations, and workspace.

Beyond storage savings, quantization also accelerates computation through hardware parallelism. The speedup emerges from how modern processors pack more operations into the same hardware resources when working with smaller data types. A register-packing estimate makes that effect concrete.

Napkin Math 10.2: The SIMD multiplier

Problem: A compute-bound layer runs INT8 instead of FP32 on the same processor. Where does the speedup come from?

Mechanism: Single instruction, multiple data (SIMD). A CPU or GPU core processes data in fixed-width vector registers (for example, AVX-512 is 512 bits wide).

Math:

1. **Register width:** 512 bits.
2. **FP32 capacity:** $512/32 = 16$ elements per vector instruction.
3. **INT8 capacity:** $512/8 = 64$ elements per vector instruction.

Result: Switching to INT8 packs 4× more elements into the same register. Throughput Gain = INT8 elements/inst/FP32 elements/inst = $64/16 = 4\times$

Systems insight: Quantization delivers up to 4× speedup on compute-bound layers from vector packing alone, on hardware whose INT8 vector instructions have comparable throughput to FP32, even before considering memory bandwidth savings.

Reducing numerical precision introduces trade-offs, however. Lower-precision formats can cause numerical instability and quantization noise, potentially affecting model accuracy. Figure 10.16 shows an illustrative residual distribution aggregated across values with heterogeneous scales and possible clipping. For an ideal uniform quantizer, within-range rounding error is instead bounded by $[-\Delta/2, \Delta/2]$; errors outside that interval arise from clipping or from aggregating different step sizes. Some architectures, such as large transformer-based NLP models, tolerate quantization well, whereas others may experience significant degradation. Selecting the appropriate numerical precision therefore requires balancing accuracy constraints, hardware support, and efficiency gains.




Figure 10.16: Illustrative Aggregate Quantization Residuals: Illustrative aggregate residual distribution across heterogeneous scales and clipped values. Within one unclipped uniform quantizer, rounding error is bounded by $[-\Delta/2, \Delta/2]$; the displayed tails represent clipping or aggregation across different step sizes.

To appreciate how precision loss manifests in practice, examine the representative residual distribution in figure 10.16: most values quantize with small error, while clipping or heterogeneous scale choices can create larger residuals that influence model accuracy. Understanding this noise is essential, but practitioners ultimately care about end-to-end speedup, and the magnitude of the quantization speedup depends on whether a workload is compute bound or memory bound.

Napkin Math 10.3: The quantization speedup (compute bound)

Problem: A compute-bound matrix multiplication (for example, in a transformer multilayer perceptron block) switches from FP16 to INT8. What is the expected speedup?

Math: On modern hardware with dedicated INT8 units:

1. **Integer throughput path:** The reference A100-class accelerator rates its dedicated INT8 matrix units at 624 TOPS against 312 TFLOP/s on the FP16 path (NVIDIA Corporation 2020a; Choquette et al. 2021). The peak throughput increase is that spec ratio: $624 \text{ TOPS} \div 312 \text{ TFLOP/s} \approx 2\times$.
2. **Memory bandwidth:** INT8 weights are half the size, so loading them from memory takes half the time.
3. **Combined effect:** For compute-bound operations, the speedup is primarily from compute throughput: $\sim 2\times$ speedup.

Systems insight: The speedup from quantization depends on the bottleneck. Compute-bound operations (large batch sizes, high arithmetic intensity I) see $\sim 2\times$ from faster INT8 units, where the gain comes from the integer matrix path rather than reduced memory traffic. The bandwidth-bound case inverts this: halving the bytes moved (FP16 to INT8) yields up to $2\times$, and larger bit-width reductions scale the gain further, as the next worked example traces.

The complementary case is bandwidth bound, where the speedup tracks the reduction in bytes moved per token rather than peak arithmetic throughput.

The result is a bandwidth-driven speedup: fewer bits moved per token means more tokens/s.

Napkin Math 10.4: The quantization speedup

Problem: A deployment scenario calls for running a 7B-parameter LLM on a device with 16 GB RAM. The weights are FP16 (2 bytes).


Math:

1. **Model size:** $7 \times 10^9 \times 2 \text{ bytes} = 14 \text{ GB}$.
2. **KV cache:** A 4,096-token FP16 KV cache uses two tensors (keys and values) across 32 layers, 32 KV heads, 128 dimensions per head, and 2 bytes per value, requiring approximately 2.1 GB.
3. **Total memory:** $14 \text{ GB} + 2.1 \text{ GB} = 16.1 \text{ GB}$. This exceeds the device capacity before OS and workspace memory.
4. **Bandwidth cost:** In this simplified full-cache-traffic model, loading weights plus the KV cache at 50 GB/s takes 323 ms per token. That is 3.1 tokens/s, too slow for chat.

Fix (INT4):

1. **Quantization:** Convert weights to INT4 (0.5 bytes).
2. **New size:** $7 \times 10^9 \times 0.5 \text{ bytes} = 3.5 \text{ GB}$ of weights, or 5.6 GB including the unchanged FP16 KV cache.
3. **New speed:** Loading that total takes 113 ms. Speed rises to 9 tokens/s.

Systems insight: Weight quantization makes this configuration fit and yields a $2.9\times$ speedup in the stated traffic model, below $4\times$ because FP16 KV-cache traffic is unchanged.



Quantization speedup depends on whether memory bandwidth or compute throughput is binding.

10.4.3 Numerical format comparison

The format decision is to choose the numerical range and hardware path that preserve accuracy while reducing bytes moved, rather than to minimize bit width blindly. Table 10.9 compares commonly used numerical precision formats in machine learning, each exhibiting distinct trade-offs in storage efficiency, computational speed, and energy consumption. Formats such as FP8 and TF32 further optimize performance, especially on AI accelerators.

FP16 and BF16 formats can provide moderate efficiency gains, while accuracy preservation depends on the model, task, and training method. Many AI accelerators, such as NVIDIA Tensor Cores and TPUs, include dedicated support for FP16 computations, enabling faster matrix operations than FP32. BF16 retains FP32's 8-bit exponent but has a 7-bit mantissa, preserving a similar dynamic range while reducing precision. FP16 has a maximum finite value of 65,504, a smallest positive normal value of approximately 6.1×10^{-5} , and subnormal values down to approximately 6.0×10^{-8} . FP16 training commonly uses loss scaling and FP32 accumulation to manage underflow and overflow, while BF16's wider exponent range reduces that risk.

Table 10.9: Numerical Precision Formats: Storage ratios reflect raw value width; compute speed and power effects depend on hardware support and workload behavior.

Precision Format	Bit-Width	Storage Reduction (vs. FP32)	Compute Speed (vs. FP32)	Power Effect	Use Cases
FP32	32-bit	Baseline (1×)	Baseline (1×)	Baseline	Training & inference (general-purpose)
FP16	16-bit	2× smaller	Up to 2× peak on selected supported hardware	Hardware- and workload-dependent	Accelerated training, inference (NVIDIA Tensor Cores, TPUs)
BF16 (Brain Floating Point)	16-bit	2× smaller	Hardware- and workload-dependent	Hardware- and workload-dependent	Training on TPUs, transformer-based models
TF32 (TensorFloat-32)	19-bit	None (stored as FP32)	Up to 8× peak on NVIDIA Ampere Tensor Cores	Hardware- and workload-dependent	Training on NVIDIA GPUs
FP8 (Floating-Point 8-bit)	8-bit	4× smaller	Hardware- and workload-dependent	Hardware- and workload-dependent	Efficient training/inference (H100, AI accelerators)
INT8 (8-bit Integer)	8-bit	4× smaller	Hardware- and workload-dependent	Hardware- and workload-dependent	Quantized inference (Edge AI, mobile AI, NPUs)
INT4 (4-bit Integer)	4-bit	8× smaller	Hardware- and workload-dependent	Hardware- and workload-dependent	Ultra-low-power AI, experimental quantization
Binary/Ternary (1-bit/2-bit)	1–2-bit	16–32× smaller	Hardware- and workload-dependent	Hardware- and workload-dependent	Extreme efficiency (binary/ternary neural networks)

Compare the three bit layouts in figure 10.17 to see exactly where the bits go—and why the trade-off between precision and numerical range differs so sharply across formats.

INT8 precision offers more aggressive efficiency improvements for inference workloads. Many quantized models use INT8 for inference, reducing raw value storage by 4× relative to FP32; realized speedups depend on hardware support, kernels, and workload behavior. INT8 is widely used in mobile and embedded AI, where energy constraints are significant.

Binary and ternary networks represent the extreme end of quantization, where weights and activations are constrained to one-bit (binary) or a small ternary set. Their raw representations are compact, but realized storage and energy depend on packing and hardware support, and task quality becomes increasingly method dependent. Our keyword-spotting lighthouse (section 6.3.5) illustrates a regime in which extreme compression may be needed to fit the ~512 KB SRAM device budget. INT8 is one candidate starting point; INT4, binary, or ternary representations spend additional numerical capacity to pursue still tighter memory and power constraints.




Figure 10.17: Floating-Point Bit Layouts: Bit allocation trades dynamic range against significand precision. FP32 uses 8 exponent and 23 fraction bits. BF16 retains FP32’s 8-bit exponent and therefore a similar numerical range, but reduces the fraction to 7 bits. FP16 uses 5 exponent and 10 fraction bits, providing more significand precision than BF16 but a narrower range that can require loss scaling during training.

The energy analysis completes the motivation for quantization: fewer bits reduce memory movement and lower arithmetic energy, but only when the target processor exposes the corresponding low-precision units. A full INT8 multiply-accumulate, for example, uses roughly 20× less energy than its FP32 equivalent. Accelerators with Tensor Cores, FP8/INT8 units, TPUs, or NPUs can compound arithmetic savings with lower memory traffic, while general-purpose CPUs without efficient low-precision paths may lose much of the benefit to packing, dequantization, or scalar fallback. The practical implication is that precision is a model-hardware decision, not a software flag.

10.4.4 Precision reduction strategies

With the hardware condition established, the question becomes *how* to reduce precision without destroying model accuracy. Naive quantization introduces errors that degrade predictions, so practitioners need structured strategies that control *where* and *how* precision is reduced.

Three approaches form a complexity ladder. Post-training quantization (PTQ) reduces precision after training, requiring no retraining and minimal engineering effort. Quantization-aware training (QAT) incorporates quantization effects into the training loop, enabling models to adapt to lower precision and retain higher accuracy. Mixed-precision training assigns different precision levels to different operations, matching precision to each layer’s sensitivity. Figure 10.18 maps quantization techniques into three progressive tiers based on implementation complexity, resource requirements, and target use cases.




Figure 10.18: Quantization Complexity Roadmap: The tiers order techniques by typical engineering and validation effort, not by guaranteed accuracy. PTQ provides a low-cost starting point; QAT and mixed precision add training or kernel control; INT4, binary, and ternary methods require increasingly model- and hardware-specific validation. A deployment may stop at any tier once it meets measured quality, latency, memory, and energy criteria, and a more complex tier is not automatically better.

The roadmap is a deployment-ordering device rather than a taxonomy of numerical formats. PTQ belongs first because it changes representation with minimal training cost; QAT and mixed precision move into the production tier because they require training-loop or kernel support; INT4, binary,

and ternary methods sit at the frontier because the accuracy and hardware assumptions become architecture-specific.

10.4.4.1 Post-training quantization

Post-training quantization (PTQ) reduces numerical precision after training, converting selected weights or activations from FP32 to lower-precision representations such as INT8 without full retraining (Choukroun et al. 2019). It reliably reduces raw value payload; latency and energy improve only when the exported graph maps to efficient low-precision kernels (Wu et al. 2020).

PTQ's key advantage is low computational cost: it requires no retraining and usually no labeled training set, although activation calibration typically needs a small representative calibration dataset. However, reducing precision introduces quantization error that can degrade accuracy, especially for tasks requiring fine-grained numerical precision. Machine learning frameworks such as TensorFlow Lite, the Open Neural Network Exchange runtime, and PyTorch provide built-in PTQ support.

The core mechanism of PTQ is **uniform quantization**, which maps floating-point values to discrete integer levels using a consistent scaling factor. Because the interval between each quantized value is constant, uniform quantization simplifies implementation and enables efficient hardware execution. For the symmetric form shown here, s is chosen from the maximum absolute value so that the resulting integers stay within the target range. The quantized value q is computed as:

$$q = \text{round}\left(\frac{x}{s}\right)$$

where:

- q is the quantized integer representation,
- x is the original floating-point value,
- s is a scaling factor that maps the floating-point range to the available integer range.

Listing 10.2 demonstrates uniform quantization from FP32 to INT8, reducing the raw value payload from 32 to 8 bits per weight while measuring the resulting quantization error; scale metadata and packing overhead are excluded. Once a model is quantized, the runtime can use integer arithmetic where supported (Gholami et al. 2022).

Listing 10.2: Uniform Quantization: Converts FP32 weights to INT8, reducing raw per-weight payload from 32 to 8 bits while measuring quantization error; scale metadata and packing overhead are excluded.

```
import torch

# Original FP32 weights
weights_fp32 = torch.tensor(
    [0.127, -0.084, 0.392, -0.203], dtype=torch.float32
)
print(f"Original FP32: {weights_fp32}")
print(f"Memory per weight: 32 bits")

# Simple uniform quantization to INT8 (-128 to 127)
# Step 1: Find scale factor
max_val = weights_fp32.abs().max()
scale = max_val / 127 # 127 is max positive INT8 value

# Step 2: Quantize using our formula q = round(x/s)
weights_int8 = torch.round(weights_fp32 / scale).to(torch.int8)
print(f"Quantized INT8: {weights_int8}")
print(f"Memory per weight: 8 bits (reduced from 32)")

# Step 3: Dequantize to verify
weights_dequantized = weights_int8.float() * scale
print(f"Dequantized: {weights_dequantized}")
print(
    f"Quantization error: "
    f"{(weights_fp32 - weights_dequantized).abs().mean():.6f}"
)
```


An alternative, nonuniform quantization, assigns finer-grained precision to numerical ranges that are more densely populated, which can preserve accuracy for models whose weight distributions concentrate around specific values. Nonuniform schemes require more complex calibration and are less common in production, but they can be effective for models particularly sensitive to precision changes.

PTQ is a useful first experiment for many vision, language, and speech models, but sensitivity varies by architecture, layer, and distribution. When a uniform PTQ configuration misses the quality target, the next options include finer granularity, mixed precision, weight-only methods, better calibration, or QAT.

Calibration An important aspect of PTQ is the calibration step, which estimates a clipping range $[\alpha, \beta]$ for quantized tensors. A poor range either clips important values or wastes integer levels on outliers. Calibration reduces that error on the sampled distribution; it does not guarantee performance on deployment data.

Post-training quantization relies on representative calibration data to estimate tensor clipping boundaries. Follow the step-by-step pipeline in figure 10.19 and the layer-by-layer procedure in algorithm 10.1, observing how activation observers produce scale metadata for quantized runtime execution.

Algorithm 10.1 Post-training quantization calibration

Require: pretrained model f_θ ; representative calibration set C ; bit width b ; granularity (layer/group/channel); method (max, entropy, percentile)

Ensure: quantized weights and per-group range metadata $\{(\alpha_g, \beta_g, s_g, z_g)\}$

- 1: attach observers to weights and selected activation tensors, at the chosen granularity
 - 2: **for** each batch in C **do** $\triangleright$ *calibration pass*
 - 3: run f_θ in inference; record per-group value distributions
 - 4: **end for**
 - 5: **for** each quantization group g **do**
 - 6: select clipping range $[\alpha_g, \beta_g]$ by the chosen method
 - 7: derive scale s_g , zero-point z_g mapping $[\alpha_g, \beta_g]$ into the b -bit integer range
 - 8: **end for**
 - 9: quantize static weights; export activation range metadata for inference
 - 10: if accuracy falls below the production threshold, widen ranges, change granularity, or move to QAT
-




Figure 10.19: Post-Training Quantization Pipeline: PTQ converts a trained FP32 model directly to INT8 without retraining. Representative calibration data passes through the model to record activation ranges $[\alpha, \beta]$, computing scale factor S and zero-point Z parameters to quantize parameters before deployment.

The single calibration pass over C avoids retraining, and FP32-to-INT8 weight quantization cuts raw stored weight values by $4\times$ before scale, zero-point, and packing metadata. The price is that the static activation ranges fixed in the loop risk saturation or wasted integer levels when the calibration data misses deployment tails. The quantization step then converts model parameters to the lower-precision format, producing the final quantized model.

For example, consider quantizing activations to 8-bit integers. Mapping an unnecessarily wide real-valued clipping interval to $-128, \dots, 127$ wastes resolution. Calibration passes a representative dataset through the model and selects the real-valued clipping interval and scale.

Common calibration methods include **max** (uses the maximum absolute value and is susceptible to outliers), **entropy** (selects a range by minimizing a KL-divergence objective), and **percentile** (clips a chosen tail fraction). Figure 10.20 shows why outlier handling matters: long activation tails can stretch the range and leave too few levels for the bulk of the distribution.




Figure 10.20: Activation Distribution: ResNet-50 layer activations exhibit a long tail, with outlier values that can lead to inefficient precision use if not handled carefully. Source: (Wu et al. 2020).

Calibration ranges can be **symmetric** (a zero-centered range, typically with zero-point zero) or **asymmetric** (one scale with a generally nonzero zero-point, useful when distributions are skewed). The choice of method and range significantly affects quantized model accuracy.

Tuning quantization ranges A key challenge in post-training quantization is selecting the appropriate calibration range $[\alpha, \beta]$ to map floating-point values into a lower-precision representation. The choice of this range directly affects the quantization error and, consequently, the accuracy of the quantized model. Figure 10.21 contrasts the two primary calibration strategies: symmetric calibration and asymmetric calibration.




Figure 10.21: Calibration Range Selection: Symmetric vs. asymmetric affine quantization mappings. Symmetric quantization (left) forces real zero to map to integer zero, constraining range selection to $[-\alpha, \alpha]$ at the expense of wasted dynamic range for skewed activation distributions; asymmetric quantization (right) uses an explicit zero-point offset Z to map non-centered ranges $[\alpha, \beta]$ efficiently.

Quantization scales can be assigned symmetrically around zero or dynamically offset for skewed activation distributions. Compare the two mapping diagrams side by side in figure 10.21, observing zero-point preservation on the left versus asymmetric range utilization on the right.

Granularity After determining the clipping range, the next optimization step is adjusting the granularity of that range to retain as much accuracy as possible. In CNNs, each convolutional filter's weights may have a different range of values. The quantization process must account for these differences to preserve model performance.

Napkin Math 10.5: Calculating scale and zero-point

The affine quantization formula: An affine map sends a floating-point range $[\alpha, \beta]$ that includes zero to an integer range $[0, 2^b - 1]$ (for example, UINT8) via the following construction.

Analysis: We need a linear mapping $x \approx s(x_q - z)$, where s is the **scale** (step size) and z is the **zero-point** (integer value corresponding to real zero). The affine quantization process consists of three steps:

1. **Calculate scale** (s): Divide the real range by the integer range using equation 10.1.

$$s = \frac{\beta - \alpha}{2^b - 1} \quad (10.1)$$

2. **Calculate zero-point** (z): Shift the range so that real zero maps to an integer using equation 10.2.

$$z = \text{round}\left(\frac{-\alpha}{s}\right) \quad (10.2)$$

3. **Quantize** ($x \rightarrow x_q$) using equation 10.3:

$$x_q = \text{clamp}\left(\text{round}\left(\frac{x}{s} + z\right), 0, 2^b - 1\right) \quad (10.3)$$

Scenario: Suppose the activations range from $\alpha = -1$ to $\beta = 3$, and the target precision is UINT8 ($b = 8$).

- **Range:** $\beta - \alpha = 4$.
- **Steps:** $2^8 - 1 = 255$.
- **Scale:** $s = \frac{4}{255} \approx 0.0157$.
- **Zero-point:** $z = \text{round}(-(\alpha)/s) = \text{round}(-(-1)/\frac{4}{255}) = \text{round}(63.75) = 64$.

Systems insight: The real value 0.0 is represented by the integer 64, which ensures that zero-padding (common in CNNs) is represented exactly, preventing “quantization drift” where padding introduces nonzero noise.

Sharing a single quantization scale across an entire layer can increase clipping or rounding error when channel ranges differ. Observe the filter-range variation in figure 10.22, contrasting shared layerwise limits with independent per-channel scales.

Quantization granularity trades calibration and metadata overhead against representation error. More local ranges can preserve information when channels differ, but they do not improve accuracy automatically and may lack an efficient kernel path. Each additional scale also has to be stored, loaded, and applied at the granularity supported by the backend. Table 10.10 summarizes four common levels, from one shared layer range to local ranges within a filter.

Table 10.10: Quantization Granularity Levels: Granularity trades calibration overhead for accuracy, from a single per-layer range to per-filter sub-ranges.

Level	Range sharing	Trade-off
Layerwise	One range per layer	Simple but suboptimal when filter ranges vary widely
Groupwise	Filters grouped with shared ranges	Used in Q-BERT (Shen et al. 2020) for transformer attention layers
Channelwise	One range per filter	Common default; balances quantization error against scale metadata and implementation overhead
Sub-channelwise	Ranges within each filter	More local ranges can reduce quantization error but increase metadata and implementation overhead

Channelwise quantization often improves accuracy over layerwise quantization when channel ranges differ, at the cost of additional scale metadata and implementation support. With granularity




Figure 10.22: Quantization Range Variation: Different convolutional filters can have different weight ranges, motivating per-filter quantization to reduce quantization error. The differing clipping scales illustrate the choice between one layerwise range and separate channelwise ranges. Source: (Gholami et al. 2022).

determined, the next consideration is what to quantize. Neural networks contain two primary numerical components: the static weights learned during training and the dynamic activations computed during inference. Each presents distinct quantization challenges.

Weights vs. activations Executing low-precision inference requires transforming floating-point inputs before kernel execution and requantizing intermediate accumulator outputs. Follow the color-coded pipeline stages in figure 10.23, observing how violet INT8 conversion blocks feed into integer SIMD matrix multiplication before red 32-bit accumulation.




Figure 10.23: Quantized Matrix Multiplication Pipeline: Color-coded computation stage. Floating-point weights and activation inputs are quantized to INT8 (violet), fed into integer SIMD/Tensor Core matrix multiplication (blue), accumulated in 32-bit registers (red), and requantized back to INT8 (green) before non-linear activation. From the HarvardX TinyML course.

Activation quantization maps layer outputs to a lower-precision representation during inference. It reduces intermediate payloads and can reduce arithmetic cost on hardware with an efficient integer path, but it also introduces error between layers. For example, a CNN may convert FP32 feature maps to INT8 before later convolutions. The benefit is greatest when consecutive operators remain on the integer path; isolated quantized operators may pay conversion cost without reducing enough work. Whether activation quantization improves latency therefore depends on supported kernels, scale handling, tensor lifetimes, and any quantize or dequantize boundaries that remain in the exported graph.

Activation-aware methods such as **activation-aware weight quantization (AWQ)**¹⁸ target weight traffic in LLM inference. This approach is relevant to the GPT-2/Llama lighthouse when decoding is limited by loading model weights. By protecting a small fraction of salient weight channels

18 | **Activation-aware weight quantization (AWQ):**

Saliency is determined by activation magnitude, not weight magnitude—a distinction that matters because a small weight multiplied by a large activation produces a large output contribution. AWQ protects about 1 percent of salient weight channels through activation-aware scaling while quantizing most weights to low-bit formats, so a 7-billion-parameter FP16 weight set that would occupy about 14 GB can move toward an INT4-class weight footprint of about 3.5 GB before metadata, scales, and kernel-specific packing overheads (Lin et al. 2024).

based on activation magnitude, AWQ supports INT4-class weight quantization while controlling task degradation. Packed low-bit weights can reduce the memory traffic required during token generation when the runtime provides a matching kernel (Lin et al. 2024).

Static vs. dynamic quantization After determining the type and granularity of the clipping range, practitioners must decide *when* the clipping ranges are calculated. Two primary approaches exist for quantizing activations: static quantization and dynamic quantization.

In **static quantization**, the clipping range is precalculated and remains fixed during inference. This method introduces no runtime range-estimation overhead. Quantize, requantize, or dequantize operations may still remain. The fixed range can, however, lead to lower accuracy compared to dynamic quantization. A typical implementation involves running calibration inputs to compute the typical activation range (Jacob et al. 2018; Yao et al. 2021).

Dynamic quantization instead calculates a range at runtime for selected tensors. The range can adapt to the current input rather than rely on one calibration distribution, which may reduce clipping error. The trade-off is runtime work for range estimation and conversion; the net quality and latency effects depend on the operators and backend.

These timing and granularity decisions interact with the broader choice of quantization methodology. Table 10.11 compares post-training quantization, quantization-aware training, and dynamic quantization, each offering distinct strengths and trade-offs for different deployment scenarios.

Table 10.11: Quantization Method Comparison: The choice turns on three binding questions, namely engineering budget, accuracy threshold, and input variability, rather than on a feature checklist.

Method	Engineering effort	Accuracy preservation	Input adaptability	Reach for it when
Post-training quantization (PTQ)	Low: no retraining	Model-dependent; parameters do not adapt	Fixed calibration range	a calibrated model meets the measured accuracy and performance targets
Quantization-aware training (QAT)	High: additional training with quantization simulation	Often stronger when PTQ falls short	Fixed calibration range	PTQ misses the accuracy target and the training budget allows
Dynamic quantization	Moderate: ranges recomputed at runtime	Model-dependent; range fits each input	Per-input range	activation ranges vary widely across inputs and runtime overhead is acceptable

PTQ is the low-effort starting point; QAT spends training budget when PTQ misses the accuracy target, while dynamic quantization spends runtime work to track input-dependent ranges. Before proceeding to advanced techniques, a quick checkpoint tests comprehension of these core quantization modes.



Checkpoint 10.3: Calibration and range choices

Quantization succeeds only when the range policy matches deployment data and runtime constraints.

Range design

- ☐ **Calibration data:** explain how an unrepresentative calibration set creates clipping or wastes integer levels.
- ☐ **Granularity:** compare per-tensor and per-channel ranges by the error and metadata trade-off each creates.
- ☐ **Static vs. dynamic ranges:** select a range policy when activation distributions shift across inputs and runtime overhead is constrained.

System estimate

- ☐ **Quantization payoff:** for a 7B-parameter model served from HBM, estimate FP16 versus INT4 weight memory and the rough decode-throughput gain for memory-bandwidth-bound inference.

PTQ in practice The preceding subsections reveal PTQ’s core trade-off: simplicity vs. accuracy control. PTQ requires no retraining, making it a common starting point for deployment optimization. Whether calibration is sufficient depends on the model, data, quantizer, target hardware, and measured accuracy requirement.

The limitation is that PTQ does not update model parameters to recover from accuracy loss. If the quantized model falls below the production threshold, practitioners can recalibrate, change granularity or quantizers, use mixed precision, or apply quantization-aware training. QAT integrates precision constraints directly into the training process using a **straight-through estimator (STE)** to pass gradients through non-differentiable quantization operators. This allows weights to adapt to low-bit representations during retraining, recovering accuracy at the cost of incurring full training compute (O_{train}) and hyperparameter tuning.

10.4.4.2 Quantization-aware training

Quantization-aware training incorporates numerical rounding noise directly into the optimization graph. Follow the training pipeline in figure 10.24, tracing how fake-quantization nodes simulate low-precision arithmetic during forward execution while straight-through estimators preserve backward gradient updates.




Figure 10.24: Quantization-Aware Training Pipeline: QAT inserts fake-quantization nodes into the forward graph to simulate low-precision clipping and rounding errors during training. Backpropagation updates floating-point master weights using straight-through estimators (STE), allowing network parameters to adapt to numerical noise before exporting fixed integer scales for inference.

QAT can reuse calibration information from a PTQ pass. Trace the two-stage pipeline in figure 10.25: PTQ estimates scales and ranges from calibration data without gradient computation. QAT then starts from the pretrained weights and fine-tunes with simulated quantization, allowing weights to adapt to low-precision constraints through backpropagation. The training cost and resulting accuracy remain workload-dependent.

Training mathematics During forward propagation, weights and activations are quantized and dequantized to mimic reduced precision. Let x be a full-precision value, s the scaling factor that maps floating-point values into a lower-precision range, and q the simulated quantized value. This process is typically represented as:

$$q = \text{round} \left(\frac{x}{s} \right) \times s$$

where q represents the simulated quantized value, x denotes the full-precision weight or activation, and s is the scaling factor mapping floating-point values to lower-precision integers.

Although the forward pass uses quantized values, gradient calculations during backpropagation remain in full precision. The Straight-Through Estimator (STE) accomplishes this,¹⁹ which approximates the gradient of the quantized function by treating the rounding operation as if it had a derivative of one. In effect, the STE pretends quantization is the identity function during backpropagation, allowing gradients to flow unchanged through otherwise nondifferentiable operations. This approach prevents the gradient from being obstructed due to the nondifferentiable nature of the quantization operation, thereby allowing effective model training (Bengio, Léonard, et al. 2013).

¹⁹ | **Straight-through estimator (STE):** Proposed by Bengio, Léonard, et al. (2013), the STE substitutes the identity function for the true gradient of rounding, which is zero almost everywhere because rounding is piecewise constant. The common identity STE copies the upstream gradient through the quantizer. It is a biased heuristic rather than the true derivative.




Figure 10.25: PTQ-to-QAT Pipeline: PTQ calibration can initialize ranges and scales before QAT fine-tunes the pretrained weights with simulated quantization. Calibration data feeds the PTQ stage, and training data feeds QAT.

Integrating quantization effects during training lets the model adapt its weights and activation ranges to simulated low-precision error. QAT can recover quality that a post-training route loses, particularly at aggressive bit widths, but the advantage depends on the model, quantizer, calibration, and training procedure (Krishnamoorthi 2018).

Fake quantization nodes and implementation QAT implementation relies on fake quantization operations that simulate quantization during forward propagation while maintaining full precision for gradient computation. These operations insert quantize-dequantize pairs into the computational graph, creating a training-time simulation of inference-time behavior.

A fake quantization node must preserve the same errors the inference runtime will see, so its three operations model the deployment path during training:

1. **Quantization:** Map floating-point value to discrete quantization level
2. **Clipping:** Enforce range constraints based on bit width
3. **Dequantization:** Convert back to floating-point for subsequent operations

Mathematically, for symmetric quantization with bit width b , given a floating-point input value x :

$$q_{\text{level}} = \text{clip} \left(\text{round} \left(\frac{x}{s} \right), -2^{b-1}, 2^{b-1} - 1 \right)$$

$$x_{\text{fake}} = q_{\text{level}} \times s$$

where $s = \frac{\max(|x|)}{2^{b-1}-1}$ is the scale factor computed from the input distribution, and x_{fake} represents the fake-quantized output that mimics INT8 values but remains in floating-point format.

For asymmetric quantization supporting unsigned integers, assume the observed range has been extended to include zero:

$$s = \frac{\max(x) - \min(x)}{2^b - 1}$$

$$z = \text{round} \left(-\frac{\min(x)}{s} \right)$$

$$q_{\text{level}} = \text{clip} \left(\text{round} \left(\frac{x}{s} + z \right), 0, 2^b - 1 \right)$$

$$x_{\text{fake}} = (q_{\text{level}} - z) \times s$$

where z is the zero-point offset enabling asymmetric range representation.

Listing 10.3 shows the QAT forward path: fake quantization is applied to inputs and weights before convolution, simulating INT8 numerics while retaining floating-point tensors for training.

Listing 10.3: QAT Convolution Forward Pass: Fake-quantizes activations and weights before convolution to simulate deployment numerics during training.

```
# Forward pass with fake quantization
def qat_conv_forward(x, weight):
    # Fake quantize input activations
    x_scale = compute_scale(x, bits=8, symmetric=False)
    x_zero = compute_zero_point(x, x_scale, bits=8)
    x_quant = fake_quantize(x, x_scale, x_zero, bits=8)

    # Fake quantize weights (typically symmetric)
    w_scale = compute_scale(weight, bits=8, symmetric=True)
    w_quant = fake_quantize(weight, w_scale, zero=0, bits=8)

    # Convolution with fake-quantized values
    output = conv2d(x_quant, w_quant)
    return output
```

The critical aspect of fake quantization is gradient handling during backpropagation. The rounding and clipping operations are nondifferentiable, requiring gradient approximation through the straight-through estimator:

$$\frac{\partial x_{\text{fake}}}{\partial x} = \begin{cases} 1 & \text{if } x \in [x_{\min}, x_{\max}] \\ 0 & \text{otherwise} \end{cases}$$

This approximation treats the quantization function as identity within the valid range, allowing gradients to flow unchanged through the fake quantization nodes except for values that exceed clipping bounds. During backpropagation, the full-precision gradient $\frac{\partial \mathcal{L}}{\partial x_{\text{fake}}}$ propagates directly to x for values within the quantization range. For values outside the range, this clipping-mask estimator blocks the gradient through the quantizer in either direction.

In practice, frameworks like PyTorch and TensorFlow implement fake quantization as custom autograd operators whose forward pass performs the quantize-dequantize round trip while the backward pass applies an STE. Scale handling depends on the method: observers may track distributions and later freeze, or scales may be learned. Batch normalization also requires deployment-consistent handling; common workflows fold or fuse it with adjacent operations and freeze its statistics before export.

QAT trade-offs QAT’s²⁰ primary advantage is additional control over quality at low precision. By incorporating simulated quantization noise during training or fine-tuning, the model can adapt to the deployed numerical path. Processors with dedicated integer units can then exploit INT8 arithmetic for faster or lower-energy inference when the exported operators are supported (Wu et al. 2020; Gholami et al. 2022). The benefit must be measured against PTQ for the same model and target, including the extra training cost and any difference in calibration, graph coverage, or kernel selection.

The cost is additional engineering complexity. QAT inserts simulated quantization into training, so teams must validate quantization schemes, scale behavior, and accuracy recovery on the target model. This overhead can make QAT less practical for very large models when training budgets are already constrained. Choukroun et al. (2019) illustrate the complementary post-training route: low-bit inference quantization for pretrained networks without full retraining.

In practice, the choice between PTQ and QAT follows a simple decision rule. Start with PTQ and measure accuracy on the validation set. If accuracy meets the production threshold, the engineering cost of QAT is not justified. If PTQ falls short, invest in QAT to recover the gap. A hybrid approach, starting with PTQ calibration and applying QAT fine-tuning only for accuracy-critical layers, often provides a useful balance.

PTQ and QAT commonly target 8-bit or 4-bit precision, but retained accuracy depends on the model, task, and quantization method. Some deployment scenarios, however, demand even more aggressive compression, pushing precision to the absolute limits of what neural networks can tolerate.

²⁰ | **Quantization-aware training (QAT):** QAT simulates low-precision behavior during training or fine-tuning so weight updates can adapt to clipping and rounding error. BERT studies report that training-aware or Hessian-aware methods can remain closer to full-precision General Language Understanding Evaluation (GLUE) baselines than simpler post-training routes in their evaluated settings (Zafir et al. 2019; Shen et al. 2020). The exact gap depends on model, task, quantizer, bit width, training budget, and evaluation distribution; a closer benchmark result is evidence for trying QAT, not a guarantee that its added training cost or exported runtime will meet a production target.

10.4.5 Extreme quantization

Extreme quantization is reserved for deployments where even INT8 or INT4 cannot meet the memory or energy budget. These techniques use binary representations with two values, often packed at one bit per value, or ternary representations with three values, which require at least two bits per value, to reduce memory use and computation (Courbariaux et al. 2015; Zhu et al. 2017). **Binarization** constrains weights and activations to two values (typically -1 and +1, or 0 and 1), reducing model size and accelerating inference on specialized hardware for binary neural networks (Rastegari et al. 2016). However, this constraint limits model expressiveness and can degrade accuracy on tasks such as image recognition or natural language processing (Hubara et al. 2018).

Ternarization extends binarization by allowing three values (-1, 0, +1), providing additional representational flexibility (Zhu et al. 2017). The zero value also enables sparsity. Both techniques require gradient approximation methods like Straight-Through Estimator (STE) to handle nondifferentiable quantization operations during training (Bengio, Léonard, et al. 2013), with QAT integration helping mitigate accuracy loss (Choi et al. 2018).

10.4.5.1 Challenges and limitations

Despite enabling ultra-low-power machine learning for embedded systems and mobile devices, binarization and ternarization face significant challenges. Performance maintenance is difficult with such drastic quantization, and speedups require optimized bit-packed kernels or dedicated hardware support (Umuroglu et al. 2017).

Accuracy loss remains a critical concern. Binary and ternary methods become candidates when higher-precision representations miss the memory or energy budget and the target offers efficient bit-packed execution (Rastegari et al. 2016; Hubara et al. 2018; Zhu et al. 2017; Umuroglu et al. 2017). Their suitability depends on the task, architecture, training method, packing format, and kernel—not on a fixed device class. A KWS-scale SRAM budget may motivate the experiment, but it does not establish in advance that binary or ternary weights preserve acceptable quality. Compare them with INT8, INT4, pruning, and architecture changes on the same workload, then measure total artifact size, latency, energy, and task quality on the target.

Before turning to architectural efficiency, the key engineering check is whether bit width, hardware arithmetic, and QAT/PTQ choice all trace back to deployment constraints.



Checkpoint 10.4: Quantization and precision checkpoint

Test your understanding of quantization before moving to architectural efficiency:

- ☐ **INT8 memory vs. energy:** can you explain why INT8 quantization provides roughly 4× memory reduction but potentially more than 4× energy reduction? Consider the energy cost of floating-point vs. integer arithmetic units.
- ☐ **QAT vs. PTQ:** do you understand the key advantage of quantization-aware training (QAT) over post-training quantization (PTQ), and when the extra training cost is justified?
- ☐ **Bit-width speedup:** can you explain why weight quantization provides nearly linear speedup with bit-width reduction for memory-bandwidth-bound LLM inference? Think about which resource is the bottleneck during autoregressive generation.

The first two optimization dimensions answer different questions: structural optimization (pruning, distillation, NAS) determines *what* to compute, and precision optimization (quantization) determines *how precisely* to represent and execute it. Their paper savings must be stated in the resource they actually change—nonzero count, FLOPs, or raw bytes—before those reductions are translated into latency or energy.

Yet practitioners often discover a gap between theory and practice. This illustrative scenario combines a 50 percent pruning factor with INT8 width reduction for a back-of-envelope 8× target, then assumes that only 1.5× is realized. That is roughly 18.8 percent of the paper target. The scenario illustrates why *theoretical compression does not imply proportional speedup*. Optimization must extend beyond the model itself to how computations execute on physical hardware.

The gap arises from several sources. Sparse matrices stored in dense format waste memory bandwidth loading zeros—the hardware cannot skip what it does not know is zero. Operations that could run in parallel execute sequentially when data dependencies or the implementation prevent overlap. A fixed-depth model gives every input the same layer budget even when some cases could be resolved earlier. Closing the gap between “optimized on paper” and “optimized in practice” is the domain of our third optimization dimension: **Architectural efficiency**. This dimension asks whether structural and precision optimizations translate into real speedups once computation patterns meet hardware capabilities.

10.5 Architectural Efficiency

Architectural efficiency starts from the execution trace. The preceding section quantified the gap between paper and measured speedup; a profiler explains where it goes: sparse tensors may still move through dense kernels, reduced-precision operators may require conversions the hardware cannot hide, and small layers may spend more time launching kernels and moving intermediates than doing arithmetic. The model has become smaller on paper, but the execution trace still asks the machine to perform an inefficient sequence of memory accesses and operations.

Where representation optimization determines *what* computations to perform and precision optimization determines *how precisely* to compute them, architectural efficiency determines *how* those computations fit the machine. Measured bottlenecks determine which architectural response is useful. Hardware-aware design changes the model before training so its layers match the deployment envelope. Sparsity exploitation makes removed weights visible to kernels that can skip them. Dynamic computation lets easy inputs leave early instead of paying for the worst case. Operator fusion reduces memory traffic when adjacent operations would otherwise write and reread the same tensors.

10.5.1 Hardware-aware design

Hardware-aware design begins before compression. If the target device cannot keep convolution kernels fed, cannot store activations without spilling, or cannot meet the power budget at the chosen input resolution, pruning and quantization only treat symptoms. The architecture itself must expose the kind of work the hardware can execute efficiently. That means choosing layer shapes, scaling rules, and operator patterns with memory bandwidth, parallelism, and energy as first-class constraints rather than post-hoc deployment checks.

10.5.1.1 Efficient design principles

The first design step is to identify what the trace shows as the limiting resource. A model can miss its target because every layer is too expensive, because one convolution family dominates arithmetic, because activations or parameters do not fit the memory hierarchy, or because the candidate architecture ignores the platform’s preferred operators. Table 10.12 organizes the common responses by the bottleneck they address rather than by model family.

Table 10.12: Hardware-Aware Design Principles: A profiling trace should determine which architectural response is appropriate. MobileNet responds to redundant convolutional computation with depthwise separable convolutions, DenseNet and SqueezeNet respond to memory pressure with feature reuse and parameter reduction, and hardware-aware NAS incorporates measured platform behavior directly into the search objective.

Observed bottleneck	Architectural response	Example networks
Over-budget model scaling	Adjust depth, width, and resolution together so the model stays within the latency, memory, and power envelope.	EfficientNet, RegNet
Redundant computation	Replace expensive dense operations with factorized or grouped operations that preserve useful channel mixing at lower arithmetic cost.	MobileNet, ResNeXt
Memory pressure	Reduce parameter and activation storage, or reuse features so the working set fits the available cache, SRAM, or device memory.	DenseNet, SqueezeNet
Platform mismatch	Include measured device latency, operator support, and power behavior in the architecture search or design loop instead of optimizing FLOPs alone.	MobileNetV3, MnasNet

These responses interact. Reducing convolutional FLOPs with depthwise separable convolutions²¹ helps only if the target runtime has efficient kernels for the resulting operators. Shrinking a model with parameter-reduction layers helps only if activation storage or memory traffic was part of the measured problem. Hardware-aware design therefore does not replace profiling; it moves profiling information earlier, into the architecture itself.

10.5.1.2 Scaling optimization

The first architectural response is global: when every stage of the profile is over budget, the problem is not one bad layer; the model is scaled incorrectly for the deployment envelope. The design task is then to distribute capacity across depth, width, and input resolution rather than choose a single parameter count. Depth increases sequential work and activation storage. Width exposes more parallel work but raises memory use. Resolution improves spatial detail while increasing the number of positions each convolution must process. The right balance depends on the machine: a highly parallel accelerator can often exploit width, while a small edge device may be dominated by memory capacity and energy per access.

Mathematically, the total FLOPs for a convolutional model can be approximated as:

$$\text{FLOPs} \propto N_L \cdot w^2 \cdot r^2,$$

where N_L is depth (number of layers), w is width, and r is the input resolution. This expression shows why naive scaling fails: increasing width and resolution together multiplies work quickly, and the resulting model may exceed the memory bandwidth or power budget even when the parameter count appears reasonable.

Compound scaling turns this balancing act into a controlled design rule. Instead of adjusting depth, width, and resolution independently, compound scaling grows all three dimensions by fixed ratios (α, β, γ) relative to a base model:

$$N_L = \alpha^\phi N_{L,0}, \quad w = \beta^\phi w_0, \quad r = \gamma^\phi r_0$$

Here, ϕ is a scaling coefficient, and α , β , and γ are scaling factors determined from empirical accuracy and efficiency measurements. The rule matters because it prevents one dimension from consuming the budget before the others can contribute useful accuracy.

EfficientNet (section 10.3.4) validated this principle by using search to find a baseline architecture and then scaling it with balanced depth, width, and resolution coefficients (Tan and Le 2019). The lesson is not that every deployment should use EfficientNet. The systems lesson is that scaling is a resource-allocation decision: the same accuracy target can imply different depth-width-resolution trade-offs depending on which resource the target platform makes scarce. Later benchmarking material formalizes how to measure those trade-offs; here, the design principle is to scale the dimensions against the binding resource.

The same logic extends beyond convolutional models. Transformer layers, attention heads, sequence length, and embedding width play roles analogous to depth, width, and resolution: each increases capacity, but each stresses compute, memory bandwidth, or activation storage differently. Hardware-aware scaling keeps those dimensions tied to the measured bottleneck instead of treating model size as a single scalar.

10.5.1.3 Computation reduction

If the profile shows that a small set of convolutional operators dominates arithmetic, reducing the whole model uniformly is wasteful. The better response is to change the expensive operator. Modern efficient architectures do this by factorizing dense computations into cheaper pieces that preserve the representation needed for accuracy.

Depthwise separable convolutions, popularized by MobileNet, exemplify this approach by decomposing standard convolutions into two stages: depthwise convolution (applying separate filters to each input channel independently) and pointwise convolution (1×1 convolution mixing outputs across channels). For batch size one, unit stride, padding that preserves an $h \times w$ spatial size, and a square kernel of width k , the computational complexity of standard convolution with C_{in} input channels and C_{out} output channels is:

$$\mathcal{O}(hwC_{\text{in}}C_{\text{out}}k^2)$$

²¹ | **Depthwise separable convolutions:** This technique reduces computation by factorizing a standard convolution into separate depthwise (per-channel) and pointwise (1×1) operations. MobileNet architectures use this factorization to trade model capacity and accuracy against lower operation counts for on-device vision (Howard et al. 2017; Sandler et al. 2018).

where k is kernel size. Depthwise separable convolutions reduce this to:

$$\mathcal{O}(hwC_{\text{in}}k^2) + \mathcal{O}(hwC_{\text{in}}C_{\text{out}})$$

eliminating the k^2 factor from channel-mixing operations and often achieving 5–10× FLOP reduction. The wall-clock benefit depends on kernel support and memory behavior: a mobile runtime with optimized depthwise kernels can convert much of this arithmetic reduction into latency savings, while a poorly supported backend may expose the factorized operations as many small, memory-bound kernels.

Other factorization patterns respond to the same diagnosis. Grouped convolutions, used in ResNeXt, partition feature maps into independent groups before merging them, reducing redundant cross-channel work. Bottleneck layers, used in ResNet, apply 1×1 convolutions to reduce feature dimensionality before expensive operations. SqueezeNet uses the same 1×1 idea to reduce parameters. These techniques improve efficiency when they reduce the operation that actually dominates the trace; they provide much less benefit when memory traffic, launch overhead, or unsupported kernels become the new bottleneck.

Arithmetic reduction is only one possible response. When profiling instead identifies memory capacity or data movement as the binding resource on the target device, the architecture must address that memory bottleneck directly.

10.5.1.4 Memory optimization

When the profile points to memory rather than arithmetic, the architecture must reduce the working set or the number of expensive memory accesses. Activations, feature maps, and parameters can exceed cache, SRAM, accelerator memory, or edge-device storage even when the FLOP count is acceptable. Memory-efficient architectures therefore try to preserve useful information while storing or moving less data.

DenseNet illustrates the feature-reuse response (Huang et al. 2017). In a traditional convolutional network, each layer computes a new set of feature maps, increasing the activation footprint as the network deepens. DenseNet connects layers so later computations can reuse earlier feature maps instead of relearning similar representations. In a standard convolutional network with N_L layers, if each layer generates g new feature maps, the total number of feature maps grows linearly:

$$\mathcal{O}(N_L g)$$

DenseNet reduces parameter redundancy through feature reuse, but concatenating retained features can increase activation storage and memory traffic. The resulting working set and runtime must be measured on the target hardware.

Activation checkpointing complements feature reuse by trading computation for memory during training. For N_L comparable layers with activation footprint A_{layer} per layer, evenly spaced checkpointing reduces the saved-activation term from $\mathcal{O}(N_L A_{\text{layer}})$ to $\mathcal{O}(\sqrt{N_L} A_{\text{layer}})$ while recomputing omitted activations during backpropagation. In the compression context, checkpointing enables training of larger models within fixed memory budgets, which in turn provides more capacity for subsequent pruning or distillation to exploit.

Parameter reduction applies the same reasoning to storage. SqueezeNet uses 1×1 convolutions to reduce the number of input channels before applying standard convolutions, making the expensive layer operate on a smaller representation (Iandola et al. 2016). The number of parameters in a standard convolutional layer is:

$$\mathcal{O}(C_{\text{in}}C_{\text{out}}k^2)$$

By reducing C_{in} using 1×1 convolutions, SqueezeNet reduces parameter count, achieving the paper's AlexNet-level accuracy target with far fewer parameters than AlexNet. That trade is attractive when the deployment constraint is flash storage, model download size, or parameter bandwidth; it is less decisive when activation memory or operator overhead dominates.

Feature reuse, activation checkpointing, and parameter reduction are therefore not interchangeable recipes. Each changes a different part of the memory problem. Reused features reduce redundant representations, checkpointing reduces training-time activation storage, and 1×1 bottlenecks reduce

parameter movement. The correct choice follows from which memory term the profile shows as binding.

Beyond reducing what data must be stored, substantial efficiency gains emerge from optimizing how operations access memory. The next technique addresses this by combining multiple operations to reduce memory traffic.

10.5.1.5 Operator fusion

Consider a typical neural network layer: convolution followed by batch normalization followed by rectified linear unit (ReLU). Without fusion, each operation writes its output to GPU global memory, then the next operation reads that output back. Three memory round-trips occur for what could be computed entirely in fast on-chip registers. By fusing these operations into a single kernel, compilers and inference engines eliminate the redundant memory transactions, improving both throughput and latency on memory-bound workloads (Chen et al. 2018; NVIDIA 2024c).



Definition 10.5: Operator fusion

Operator fusion is a compiler and runtime optimization that combines adjacent tensor operations into a single fused kernel so intermediate values remain in registers or on-chip memory instead of being written to and reread from accelerator global memory.

1. **Significance:** In an ideal chain of N element-wise operations over the same M -byte tensor, separate kernels can move about $2NM$, while a fused kernel can approach $2M$ by keeping intermediates on chip. Real traffic also includes weights, caches, alignment, and spills, so profiling must confirm the reduction.
2. **Distinction:** Unlike pruning, quantization, or distillation, operator fusion does not change the model's parameters, precision, or architecture. It changes the execution schedule of mathematically equivalent operations, preserving model outputs while improving latency and throughput on memory-bound workloads.
3. **Common pitfall:** A frequent misconception is that more fusion is always better. Fusion is constrained by data dependencies, tensor shapes, register pressure, and cache capacity; over-fusing can reduce occupancy or force spills back to memory, erasing the benefit.

Modern neural networks consist of sequences of operations such as convolution, batch normalization, activation functions, and element-wise operations. When executed independently, each operation requires four steps:

1. Loading input tensors from global memory
2. Performing computation
3. Writing output tensors back to global memory
4. Launching the next kernel

The read-compute-write cycle creates memory bandwidth bottlenecks for operations with low arithmetic intensity (FLOP/byte). For an ideal sequence of N operations over the same M -byte intermediate tensor, the unfused traffic is:

$$D_{\text{vol,unfused}} = 2NM$$

where each operation reads and writes the intermediate once. If the sequence is legal to fuse and intermediates stay on chip, the traffic can approach:

$$D_{\text{vol,fused}} = 2M$$

by reading the intermediate input once and writing its final output once. Weights, auxiliary inputs, cache effects, and register spills are outside this simple bound. Several inference patterns nevertheless approximate it closely enough for fusion to matter.

Convolution-BatchNorm-ReLU fusion The common Conv-BN-ReLU pattern in listing 10.4 can replace three launches and intermediate round-trips with one fused path when the compiler can legally fold batch normalization and combine the activation.

Listing 10.4: Conv-BN-ReLU Fusion: Combining three operations into a single kernel reduces memory traffic from 6 transfers to 2, eliminating intermediate memory writes.

```
# Pseudocode: runtime APIs and fusion legality are
# implementation-dependent.
# === UNFUSED: 3 kernel launches, 6 memory transfers ===
conv_out = conv2d(input, weight)
bn_out = batch_norm(conv_out, ...)
relu_out = relu(bn_out)

# === FUSED: 1 kernel launch, 2 memory transfers ===
def conv_bn_relu_fused(input, weight, gamma, beta, mean, var):
    # Read input and weight once
    conv = conv2d(input, weight)

    # Apply batch norm in registers (no memory write)
    bn = gamma * (conv - mean) / sqrt(var + eps) + beta

    # Apply ReLU in registers (no memory write)
    output = relu(bn)

    # Write final result once
    return output
```

The arithmetic operations remain identical, but memory traffic drops from 6 transfers to 2 transfers (3× reduction). For a ResNet-50 layer with 256 channels and spatial size 28×28 , this eliminates $4 \times 256 \times 28 \times 28 \times 4$ bytes ≈ 3.2 MB of intermediate memory traffic per layer.

The same principle extends beyond CNNs. General matrix multiplication (GEMM) bias-activation fusion eliminates intermediate writes in transformer linear layers by computing element-wise operations in registers immediately after each matrix multiplication output element. Attention tiling, as in FlashAttention,²² reduces HBM traffic by processing attention in SRAM-sized tiles and avoiding materialization of the full $S \times S$ attention matrix, as detailed in section 8.6.4.

Memory bandwidth analysis quantifies these fusion benefits concretely. Consider a Conv-BN-ReLU sequence operating on a $28 \times 28 \times 256$ feature map (802.8 KB). Without fusion, each operation performs its own memory round-trip: Conv reads input (802.8 KB) plus weights (2.4 MB) and writes output (802.8 KB), totaling 4 MB. BN then reads that output, adds its parameters (2 KB), and writes again, for 1.6 MB. ReLU repeats the pattern for another 1.6 MB. The total unfused memory traffic is 7.2 MB. With fusion, the entire sequence reads input and weights once and writes the final output once, requiring only 4 MB—a 44.5 percent bandwidth reduction.

At the modeled 900 GB/s HBM bandwidth, these traffic volumes imply ideal lower bounds of 8 microseconds before fusion and 4.5 microseconds after fusion, a 1.80× ratio for the traffic component alone. These are not full layer latencies because convolution compute, cache effects, and launch overhead must also be included.

As table 10.13 shows, fusion benefits vary by workload. Memory-bound operations benefit most, while compute-bound operations see minimal improvement.

Table 10.13: Illustrative Operator Fusion Benefit by Workload: Memory-bound workloads tend to benefit more than compute-bound workloads; realized gains depend on the graph, runtime, sequence shape, and hardware.

Workload	Relative Fusion Benefit	Why
Element-wise operations	High	Highly memory bound, low arithmetic intensity
Conv-BN-Act patterns	Moderate	Mixed memory / compute characteristics
GEMM-based operations	Low	Compute bound; fusion reduces the memory-bound tail

Unfused: 6

Fused: 2

Operator fusion cuts memory transfers from six to two.

22 | FlashAttention: Demonstrates fusion’s power for memory-bound attention by tiling computation to SRAM, avoiding materialization of the full attention matrix, and reporting multi-fold speedups on long sequences (T. Dao et al. 2022). This exemplifies how operator fusion transforms memory-bound bottlenecks: the arithmetic is mathematically equivalent, but the memory access pattern changes, making longer-context attention feasible on hardware that could not otherwise afford the full intermediate matrix.

Workload	Relative Fusion Benefit	Why
Attention mechanisms	High	Long sequences amplify avoided attention-intermediate traffic

Fusion also reduces kernel launch overhead. Each CUDA kernel launch incurs microsecond-scale latency. In a hypothetical graph with fifty-three eligible Conv-BN-ReLU triplets, unfused execution launches 159 kernels, while fused execution launches 53 kernels, saving repeated launch overhead in addition to memory traffic. ResNet-50 does not contain fifty-three direct triplets because residual additions interrupt many such patterns.

Fusion implementation spans the software stack, from framework-level pattern matching through compiler optimization such as Accelerated Linear Algebra (XLA), TVM, and TensorRT to runtime fusion that adapts to input shapes and hardware characteristics. The process becomes tractable when a framework can export or trace the relevant computation as a graph whose nodes are tensor operations and whose edges record data dependencies. A compiler can then match legal patterns such as Conv→BN→ReLU and rewrite them as fused operations while checking shapes, aliases, and numerical constraints. Dynamic control flow, in-place mutation, or unsupported operations can prevent that graph rewrite, so the exported plan must be inspected rather than assumed. Section 11.8.1.3 examines the compiler and hardware dimensions of fusion in detail, including register pressure, graph matching, and platform-specific trade-offs across GPU, TPU, and edge accelerators.

Operator fusion optimizes how operations execute by reducing memory traffic between fixed computational steps. A complementary approach challenges the assumption that all computational steps must execute at all. This leads to adaptive computation methods that vary the amount of work performed based on input characteristics.

10.5.2 Adaptive computation methods

The preceding techniques (hardware-aware design and operator fusion) optimize a fixed execution graph, so absent input-dependent control, every input follows the same scheduled path. Consider image classification: a photo of a cat against a plain background may reach a confident prediction earlier than a cat partially hidden in a cluttered room. **adaptive inference** challenges the fixed-path assumption by turning inference into a control loop: measure confidence or context, select a path or stopping point, and pay the overhead of that decision. This flexibility can reduce average work when the controller reliably separates inputs that need different amounts of computation and the runtime executes those paths efficiently, but it also introduces routing, batching, calibration, and evaluation costs.

10.5.2.1 Dynamic schemes

When inputs vary in difficulty, a fixed computational budget may spend more work than some cases require. Dynamic schemes modify the inference path so average computation can fall while a full path remains available. That availability does not preserve worst-case quality automatically; the exit or routing policy must be calibrated and evaluated. The design question is what the controller may vary: depth, path, expert subnetwork, or a continuous compute budget.

The first control variable is depth. **Early exit** attaches lightweight classifiers to intermediate layers and stops once confidence is high enough (Teerapittayanon et al. 2017). BranchyNet implements this idea with multiple exit points, while multi-exit vision transformers attach the same decision rule to transformer layers. Inputs that satisfy an exit criterion can avoid later layers; the realized latency or energy saving must exceed the cost of the exit classifiers and routing.

The systems trade-off depends on where those exits run. Mobile and edge deployments may reduce both work and latency when the early classifiers are cheap (Hu et al. 2020). GPU and TPU deployments must balance any saved layers against routing overhead and batch fragmentation (Chen et al. 2024). Figure 10.26 traces the decision logic: each layer refines the representation, the confidence estimator decides whether to exit, and the remaining inputs continue.

Early exits decide how far to proceed along one path. The next control variable is route. **Conditional computation** decides which path, layer, unit, or expert should run for a given input (Bengio




Figure 10.26: Early Exit Architecture: Transformer layers dynamically adjust computation by classifying each layer's output and enabling early termination if sufficient confidence is reached, reducing average latency and power when many inputs can exit before the final layer. Source: (Xin et al. 2021).

et al. 2015). The control signal can skip work, change the computation being applied, or route representations through different structures; in every case, the model becomes an input-dependent execution graph rather than a fixed sequence of layers.

Representative mechanisms make that range concrete. SkipNet uses a lightweight gate to skip CNN layers when the input is simple and to execute the full network when the input is difficult (X. Wang et al. 2018). Dynamic Filter Networks condition the filters themselves on the input, adapting feature extraction at runtime rather than merely skipping operations (Jia et al. 2016). Capsule Networks use routing to assign lower-level capsules to higher-level capsules based on agreement; in that case, routing supports part-whole representation learning rather than guaranteeing lower operation count (Sabour et al. 2017). The systems question is therefore not just whether routing is possible, but whether the routing decision saves more time and energy than it consumes.

At subnetwork scale, the route becomes an expert-selection decision. The **mixture-of-experts (MoE)** framework uses a gating network to select a small subset of expert subnetworks rather than activating the entire model (Shazeer et al. 2017). A question about mathematics and a question about history can use different experts while sharing the same overall model capacity. Google's Switch Transformer²³ instantiates this idea in transformers by replacing the dense feedforward layer with an expert-routed layer (Fedus et al. 2022). The benefit is parameter capacity without proportional per-token compute; the cost is that the full expert pool must still stay resident in memory even though only a fraction runs per token, on top of load balancing, routing overhead, and more complicated batching. This chapter treats MoE at single-system scale; large-scale deployments add distributed expert placement across machines.

Sparse Mixture-of-Experts architectures replace monolithic feedforward networks (FFNs) with dynamically routed sub-networks. Compare the transformer block structure on the left of figure 10.27 with the expanded gating router on the right, observing how each input token activates only a single expert FFN block.

Gate-based conditional computation is effective for multi-task and transfer learning settings where inputs benefit from specialized processing pathways, but the gate is now part of the system's critical path. Efficient deployment on GPUs, TPUs, or edge devices requires scheduling and batching expert activations so that specialization does not leave accelerator lanes idle (Lepikhin et al. 2021).

Early exit and conditional computation make discrete choices: exit or continue, activate this expert or that one. Adaptive inference treats computation more like a dial, continuously modulating depth and resource allocation based on confidence and task complexity (Yang et al. 2020). Fast Neural Networks adjust the number of active layers from a real-time complexity estimate (J. Wu et al. 2019), while dynamic layer scaling progressively increases depth when uncertainty remains high. The autonomous-driving example makes the trade-off concrete: lane detection may need a shallow path, while dense multi-object tracking may need deeper processing. The system gains only if the control signal is cheap, the routed paths preserve accuracy, and the runtime can still batch enough similar work to keep hardware busy.

²³ | **Switch Transformer:** Fedus et al. (2022) scales to roughly 1.6 trillion total parameters while routing each token to one expert; in a compute-matched experiment, Switch-Base with 64 experts reached the T5-Base quality target in about one-seventh the training time; the trillion-parameter comparison reported roughly 4× speedup over T5-XXL. Routing each token to one expert reduces communication relative to top-*k* routing, but load imbalance requires auxiliary losses and capacity factors. This trade-off—massive parameter capacity at low per-token compute cost, but with complex systems engineering for load balancing—defines the MoE design space.




Figure 10.27: Switch Transformer Architecture: Replacing dense feedforward layers with a Switching FFN layer (left) enables sparse conditional routing. A gating network routes each token to a single expert (right) out of a large pool, decoupling total model capacity (parameters) from per-token compute cost (FLOPs). Source: (Fedus et al. 2022).

10.5.2.2 Implementation challenges

The efficiency gains from dynamic computation come at the price of making the control loop itself correct, cheap, and hardware-aligned. Training is harder because discrete gating decisions cannot be optimized with standard backpropagation without reinforcement learning, continuous relaxations, or regularization that stabilizes gradients across different paths. Runtime overhead then becomes the practical test: a gate that saves one layer but adds synchronization, memory traffic, or queueing delay may lose to the dense baseline. Hardware utilization compounds the problem because modern accelerators favor regular, predictable batches; when each input follows a different path, some lanes sit idle unless the runtime regroups similar paths or uses specialized kernels. Section 11.10.2 examines hardware-aware runtime strategies for these adaptive execution patterns.

The quality risks are equally important because the control policy decides which inputs receive computation. A poorly calibrated gate can underallocate work to rare but important inputs, creating biased predictions precisely where coverage matters most. If adversarial or malformed inputs can influence the gate, the same policy can become a denial-of-quality or denial-of-service lever by steering work toward cheap paths that miss hard cases or expensive paths that overload the service. Evaluation must therefore report more than average FLOPs: it must include path distributions, tail latency, accuracy by input difficulty, routing stability, and reproducibility under changing batches. Overcoming these challenges requires robust training techniques, hardware-aware execution strategies, and evaluation frameworks that account for adaptive scaling. Where dynamic computation decides whether to perform certain operations, sparsity exploitation addresses a complementary question: how to accelerate computation when many operands are zero.

10.5.3 Sparsity exploitation

Recall that pruning (from section 10.3.1) introduces zeros into weight matrices. Sparsity exploitation asks how to accelerate computation when those zeros are present. Pruning creates zeros or removes structure; sparse formats can reduce storage, and matching kernels can reduce computation. **Sparsity**²⁴ in machine learning refers to the condition where a significant portion of the elements within a tensor, such as weight matrices or activation tensors, are zero or nearly zero.

²⁴ | **Sparsity:** From Latin *sparsus* (scattered), past participle of *spargere* (to scatter); in ML, L1 regularization such as the lasso can induce exact zeros rather than merely small values (Tibshirani 1996). Software can represent arbitrary sparsity patterns, while efficient hardware execution generally requires a supported format or structured pattern (for example, NVIDIA's 2:4 sparsity). The gap between representable and executable sparsity separates theoretical compression from realized speedup.

More formally, for a sparse weight matrix $\mathbf{W}_{\text{sparse}} \in \mathbb{R}^{m \times n}$, the sparsity ratio ρ_{sparse} can be expressed as:

$$\rho_{\text{sparse}} = \frac{\|\mathbf{1}_{\{(\mathbf{W}_{\text{sparse}})_{ij}=0\}}\|_0}{m \times n}$$

where $\mathbf{1}_{\{(\mathbf{W}_{\text{sparse}})_{ij}=0\}}$ is an indicator function that yields one if entry (i, j) is zero and 0 otherwise, and $\|\cdot\|_0$ represents the L0 norm, which counts the number of nonzero elements. Exact zero is representable in floating point, but thresholded definitions also group small-magnitude values when the pruning method treats them as removable. The thresholded sparsity ratio becomes:

$$\rho_{\text{sparse}, \epsilon} = \frac{\|\mathbf{1}_{\{|(\mathbf{W}_{\text{sparse}})_{ij}| < \epsilon\}}\|_0}{m \times n}$$

where ϵ is a small threshold value.

Exact sparsity can emerge during training through regularization or be introduced deliberately through pruning, which forces selected elements to zero. It yields memory, compute, or energy savings only when the storage format and execution path skip those zeros with less overhead than dense execution. That qualification is especially important on resource-constrained devices, where index metadata and irregular accesses consume the same scarce bandwidth and power.

10.5.3.1 Sparsity types

The hardware decision begins with the pattern of zeros. Sparsity in neural networks falls into two broad categories: unstructured sparsity and structured sparsity.

Unstructured sparsity occurs when individual weights are set to zero without any specific pattern, typically through magnitude-based pruning. While highly flexible, unstructured sparsity is less efficient on hardware because it lacks a predictable structure.²⁵ Exploiting it requires specialized hardware or software optimizations.

Structured sparsity removes regular groups such as filters, neurons, channels, blocks, or fixed within-group patterns. Regularity makes storage and execution more predictable, but acceleration still depends on whether the target supports that structure and whether the resulting shapes use its compute units efficiently. It is therefore a candidate, not an automatic preference, when deployment requires predictable resource use.

10.5.3.2 Sparsity utilization methods

A sparse model with 90 percent of weights zeroed may still run at nearly full computational cost on hardware not designed for irregular memory access. The critical question is how to translate theoretical zeros into actual speedup. The processor *cannot* skip a multiplication unless it knows the operand is zero—and discovering that requires loading the operand from memory in the first place. Bridging this gap requires specialized utilization methods and hardware support that can efficiently skip zero-valued computations (Hoeffler et al. 2021). Han et al.’s pruning work (2015) is a canonical example of turning dense networks into sparse ones by removing unimportant connections, but accelerator speedups depend on whether the resulting sparsity pattern matches hardware-supported formats.

The simplest utilization method is a sparse matrix operation, which stores nonzero values and uses their indices to skip arithmetic on zeros. Consider the difference: multiplying a dense 4×4 matrix with a vector typically requires 16 multiplications, while a sparse-aware implementation computes the six nonzero products in addition to processing the sparse metadata:

$$\begin{bmatrix} 2 & 0 & 0 & 1 \\ 0 & 3 & 0 & 0 \\ 4 & 0 & 5 & 0 \\ 0 & 0 & 0 & 6 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \begin{bmatrix} 2x_1 + x_4 \\ 3x_2 \\ 4x_1 + 5x_3 \\ 6x_4 \end{bmatrix}$$

The deployment choice is therefore not a generic desire for fewer parameters; it is a choice between representations the runtime can execute efficiently. Low-rank approximation, covered earlier in

²⁵ | **Unstructured sparsity and SIMD waste:** Modern CPUs and GPUs process data in vector or tensorized groups; unstructured sparsity scatters nonzero elements irregularly through memory, so a vector load may bring back mostly zeros while still paying the full memory access cost. The processor cannot skip zero elements without first knowing where the nonzeros are, and the metadata needed to answer that question also consumes bandwidth. This is why structured sparsity can deliver speedups at lower sparsity levels than arbitrary unstructured sparsity, while unstructured sparsity often needs very high zero fractions and specialized kernels before arithmetic savings overcome indexing and lane-utilization overheads (Hoeffler et al. 2021).

section 10.3.3.1, replaces a dense matrix with smaller dense factors, while sparsity exploitation skips literal zero-valued weights. Sparsity-aware training and sparse gradient descent can help models learn or maintain zero patterns, but runtime speedup appears only when the deployed representation uses formats and kernels that skip zeros. That is why the next design question is not merely how sparse the matrix is, but what structure the zeros have.

10.5.3.3 Structured patterns

Achieving actual speedups from sparsity requires hardware that can efficiently skip zero-valued computations. Different processor architectures handle sparse patterns with varying effectiveness—for example, Ampere Sparse Tensor Cores exploit fine-grained 2:4 structured patterns while systolic arrays require dense block structures, hardware mechanics detailed in section 11.4.5. Software libraries such as cuSPARSE can help bridge this gap by reformulating sparse computations into patterns that current hardware handles efficiently. For example, MegaBlocks (Gale et al. 2022) reformulates sparse Mixture of Experts training into block-sparse operations, grouping routed expert-token work into dense tiles so specialized kernels can maintain high accelerator utilization despite irregular sparsity patterns.

A sparse pattern earns hardware speedup only when it is regular enough for kernels to predict. Two prominent formats make that regularity explicit: block sparse matrices and $N:M$ sparsity patterns. Block sparse matrices isolate blocks of zero and nonzero dense submatrices so that operations on the large sparse matrix can be re-expressed as a smaller number of dense operations on submatrices. This structure supports more efficient storage of dense submatrices while maintaining shape compatibility for matrix or vector products. For example, figure 10.28 shows how NVIDIA’s cuSPARSE (NVIDIA 2020a) library supports sparse block matrix operations and storage. Several other works, such as Monarch matrices (Tri Dao et al. 2022), have built on this block-sparsity approach to strike an improved balance between matrix expressivity and compute/memory efficiency.




Figure 10.28: Block Sparse Representation: Grouping non-zero weights into dense sub-blocks (3×3 tiles) and storing an auxiliary index array allows sparse matrices to be computed via regular sub-block GEMM kernels, restoring high GPU memory alignment and Tensor Core utilization. Adapted from NVIDIA cuSPARSE documentation (NVIDIA 2020a).

Similarly, $N:M$ sparsity retains at most N nonzeros in each group of M consecutive elements; hardware-oriented pruning often enforces exactly N retained values so every group has the same representation (Zhou et al. 2021). The deterministic format lets a supported kernel predict the value and metadata layout while retaining more capacity than removing whole channels or blocks. That regularity creates an executable compromise between arbitrary masks and dense computation, but

only for targets with the matching path. Figure 10.29 compares dense and 2:4 matrix multiplication, while STEP examines learning more general $N:M$ masks for inference (Lu et al. 2023).

At accelerator level, the same pattern-specific rule holds: hardware support is a contract between the sparse format and the execution path.





Figure 10.29: 2:4 Structured Sparsity GEMM: Left: standard dense matrix multiplication on Tensor Cores using full 8-element rows. Right: 2:4 sparse multiplication retains two nonzeros in every group of four and uses compact indices to select matching elements from the dense right-hand matrix. A supported sparse inner loop can halve the multiplied values, although metadata handling and the rest of the execution path determine end-to-end speedup. Source: PyTorch blog (PyTorch 2022).

Supported Sparse Tensor Core paths on NVIDIA Ampere and later can accelerate 2:4 sparsity by skipping prescribed zeros and carrying compact metadata, with an ideal sparse arithmetic-rate improvement of up to $2\times$ (NVIDIA Corporation 2020a). Conversely, **unstructured pruning** zeroes arbitrary individual weights. Those irregular locations can disrupt memory coalescing and lane utilization, while Compressed Sparse Row (CSR) or Coordinate (COO) metadata adds storage and non-unit-stride accesses. An unstructured sparse model can therefore run slower than its dense counterpart when metadata and underutilization outweigh skipped arithmetic. TPUs are useful contrast cases for dense systolic-array acceleration (Jouppi et al. 2021), but sparse acceleration still depends on a specific hardware and software path rather than accepting arbitrary masks. Field-programmable gate arrays can implement application-specific sparse formats, yet their efficiency depends on the chosen dataflow and resource budget; programmability is not a universal sparse-speedup guarantee.

Across all platforms, sparse operations can reduce memory bandwidth requirements and energy consumption when the sparse representation and kernels actually skip data movement rather than only arithmetic. This benefit compounds with quantization: a sparse INT8 model can require less memory traffic than either technique alone when the format overhead is small enough (Hoefler et al. 2021; Gale et al. 2020).

10.5.3.4 Challenges and limitations

The same format-hardware contract explains why sparsity often disappoints in practice. The central challenge is the gap between theoretical and practical speedups. Unstructured pruning removes individual weights based on importance, creating irregular patterns that common dense accelerator paths cannot exploit; skipping those zeros requires a sparse representation and a compatible kernel. Pruning itself adds training or analysis cost because identifying removable weights can require sophisticated importance estimation on large models. Even after sparsity is achieved, storage formats add indices whose overhead can offset saved values and arithmetic. Section C.1.5 details the compressed sparse row layout and quantifies how its per-nonzero metadata makes the memory payoff density-dependent. The performance break-even point varies with tensor shape, sparse format, kernel, batch size, and hardware; there is no universal sparsity threshold.



The accuracy-efficiency trade-off requires measurement across candidate sparsity levels. A model may tolerate one level with little measured impact and then degrade sharply after a comparatively small additional pruning step. The operating point is the highest useful sparsity whose task quality and deployed performance both satisfy their thresholds.

Energy efficiency is not guaranteed. While sparse operations reduce arithmetic operations, the overhead of sparse indexing and irregular memory access can increase power consumption on hardware not optimized for sparse patterns. On edge devices with tight power budgets, these overheads may outweigh the benefits.

Finally, sparsity benefits vary by layer, model, and execution target. A workload whose useful tensors remain dense, or a target without a matching sparse path, may see no improvement and can regress after metadata and irregular-access costs are included.

10.5.3.5 Combined optimizations

Pruning, quantization, operator fusion, dynamic computation, and sparse execution share weights, representations, and physical resources. A deployment may combine them when one technique cannot satisfy every constraint, but the combined result is not guaranteed to exceed the best individual technique. Each transformation can alter the accuracy, shapes, distributions, and bottlenecks assumed by the next, so combinations require joint quality and performance measurement rather than a promised compression ratio (Hoefler et al. 2021).

The interaction between sparsity and pruning is the most direct: pruning creates sparsity, but the pattern determines hardware efficiency. Removing entire filters or layers can produce smaller dense shapes, while fixed within-group patterns can map to supported sparse kernels. Unstructured pruning creates irregular patterns that require specialized formats and kernels to realize speedups (Elsen et al. 2020; Gale et al. 2019).

Combining sparsity with quantization can multiply raw value-payload reductions, but the deployed representation carries both sparse and quantization metadata. GPUs with dedicated sparse tensor cores can accelerate supported structured patterns, while CPU outcomes depend on the chosen format, low-precision sparse kernel, tensor shape, and indexing overhead (NVIDIA Corporation 2020a; Hoefler et al. 2021; Gale et al. 2020).

The recurring theme across all combinations is hardware alignment. Efficient model designs such as depthwise separable convolutions (Howard et al. 2017), dynamic computation, and sparsity help only when the target hardware supports the resulting operation patterns (Hoefler et al. 2021). Selecting technique combinations requires understanding target platform capabilities, as explored in chapter 11.

The coordination challenges inherent in combining sparsity with other techniques point to a broader principle: optimization techniques rarely succeed in isolation, and their effectiveness depends on sequencing decisions and hardware alignment.

Systems Perspective 10.4: The optimization composition problem

Optimization techniques interact through shared representations and physical resources rather than composing as independent functions. Pruning changes the surviving value distribution, tensor shapes, or sparse format; quantization adds scales and changes the available kernels; fusion changes which intermediates are materialized and can alter the profile that justified earlier choices. Effective optimization therefore treats the model-runtime-hardware combination as a coupled system, measures it after every transformation, and revisits the sequence when a new bottleneck appears. This makes compression a systems engineering problem as well as a machine learning problem.

10.6 Technique Selection

Knowing *how* each technique works is necessary but not sufficient; the practical question is *which* techniques to apply for a given deployment target and how to sequence them. An engineer deploying a transformer model faces a decision: the model exceeds device memory by $3\times$, inference latency is $4\times$ above the service-level objective, and the power budget allows no more than 2 W sustained.

The choice of whether to quantize first, prune first, distill to a smaller architecture, or combine techniques depends on which constraint is binding, what accuracy loss is tolerable, and how much engineering time is available. The following framework structures that decision.

Before choosing a sequence, separate the major levers by what they change. Pruning changes the operation pattern and only becomes fast when the resulting structure maps to kernels. Quantization changes bit width and is usually the first memory and bandwidth lever. Distillation changes the model itself by training a smaller dense student. Table 10.14 summarizes the first-order trade-offs; the sections that follow then refine the choice by constraint, hardware support, and engineering budget.

Table 10.14: Optimization Technique Trade-Offs: These rows identify first-order candidates rather than guaranteed outcomes. Quality risk and training cost depend on the model and method, while hardware dependency determines whether pruning or quantization changes latency. Pruning can remove structure or create sparse tensors, quantization directly reduces raw value bytes, and distillation spends teacher inference plus a student training run to produce an ordinary dense model. The candidate must still be evaluated against the same task and deployed-system thresholds.

Technique	Primary Goal	Quality Risk	Training Cost	Hardware Dependency	Candidate Use
Pruning	Reduce FLOPs/size	Pattern-dependent	Often includes fine-tuning	High for sparse speedup	Removable structures/weights
Quantization	Reduce bytes/precision	Bit-width-dependent	Low for PTQ; higher for QAT	High for runtime speed	Memory- or bandwidth-bound path
Distillation	Train smaller model	Student-dependent	Teacher and student run	Low for dense execution	A new training run is feasible

10.6.1 Mapping constraints to techniques

The binding constraint should determine the optimization family before a team chooses a specific technique. Table 10.15 connects each system constraint to the representation, precision, or architectural dimension most likely to relieve it, so technique selection starts from the deployment bottleneck rather than from the most familiar compression method.

Table 10.15: Optimization Dimensions: System constraints drive optimization along three core dimensions: model representation, numerical precision, and architectural efficiency, each addressing different resource limitations and performance goals. ✓ marks a direct lever, while △ marks a workload- or hardware-dependent benefit; for example, low-precision arithmetic helps only when the accelerator and runtime provide a supported path.

System Constraint	Model Representation	Numerical Precision	Architectural Efficiency
Computational Cost	✓	△	✓
Memory and Storage	✓	✓	△
Latency and Throughput	✓	△	✓
Energy Efficiency	✓	✓	✓
Scalability	✓	△	✓

Although each system constraint primarily aligns with one or more optimization dimensions, the relationships are not strictly one-to-one. Many optimization techniques affect multiple constraints simultaneously. Structuring model optimization along these three dimensions allows practitioners to analyze trade-offs more effectively and select optimizations that best align with deployment requirements.

10.6.2 Decision framework

The binding constraint of the deployment target determines which technique to reach for first, because each optimization addresses a different resource bottleneck. When model size is the primary constraint, as with over-the-air updates or storage-limited devices, quantization provides the most direct raw-payload reduction. Relative to FP32 storage, INT8 post-training quantization cuts raw weight bytes by 4× and requires no retraining; its quality impact is model-dependent. INT4 cuts

raw payload by $8\times$, with quality likewise dependent on the model, task, and method. When that loss is unacceptable, distillation to a smaller architecture followed by quantization is a candidate, but both the student and final low-precision artifact must meet the task threshold.

When latency is compute-bound, the optimization must reduce executed work or improve the rate at which the target performs it. Structured pruning can remove channels or filters and produce smaller dense operations, but awkward dimensions or a new memory bottleneck can limit wall-clock gains. If the target supports INT8 execution, quantization can improve arithmetic throughput as well. Early-exit architectures add a conditional option when input difficulty varies enough to repay classifier, routing, and batching overhead.

Small-batch LLM generation often presents a distinct bottleneck: autoregressive decoding can be dominated by *memory bandwidth* rather than *compute*, because each token generation loads the weight matrices but performs relatively little arithmetic. Weight-only quantization (INT4 or INT8 weights with FP16 activations) can then approach speedup proportional to the reduction in weight bytes until KV-cache traffic, runtime overhead, or larger batches shift the bottleneck.

When energy or power binds, quantization is a candidate because it can reduce arithmetic and memory energy, while structured pruning can reduce operation count. Their gains may compound when pruned shapes, low-precision kernels, and the memory path remain efficient; end-to-end measurement must also include control, memory, sensors, communication, and idle power.

These choices also depend on the available engineering budget. If PTQ misses the quality target and fine-tuning is feasible, QAT may recover part of the gap at the same precision. Distillation becomes attractive when a dense student is worth teacher inference and a separate training run, and custom NAS when unusual constraints and deployment scale justify search. Rapid experiments should start from measured baselines and reversible transformations such as calibrated PTQ, then add training-aware methods only when the observed gap warrants their cost rather than assume fixed hour-or-day timelines.

This decision framework provides starting points for individual technique selection. Validating that a chosen technique actually achieves its intended goal requires systematic profiling and measurement, which section 10.8 formalizes in detail. However, production deployments rarely rely on a single technique. Combining pruning with quantization, or distillation with hardware-aware design, introduces interaction effects that can either amplify benefits or create unexpected accuracy degradation. The following section addresses how to sequence and combine techniques effectively.

10.7 Optimization Strategies

The chapter's illustrative BERT mobile pipeline reduces a BERT-Base-sized footprint from 440 MB to 28 MB, a $16\times$ raw-footprint ratio, through staged structural change, distillation, and quantization. This is an accounting and sequencing scenario rather than one reported deployment. Each stage changes a different object: pruning changes structure, distillation trains a smaller model against teacher behavior, and quantization changes numerical representation. Their combined quality and runtime still require measurement.

Example 10.3: BERT-Base mobile deployment pipeline

Scenario: A modeled deployment starts from a 110M-parameter BERT-Base-sized language model for a resource-constrained mobile device; DistilBERT provides related evidence for training smaller transformer students (Sanh et al. 2019).

Diagnosis: Structural pruning, student training, and quantization interact through tensor shapes and learned distributions. Applying them independently or in a fixed assumed order can miss the final quality or runtime target.

Systems lesson: In this modeled pipeline, structural reduction, distillation, and QAT produce a $16\times$ footprint ratio (440 MB to 28 MB) with 0.6 percent task-quality loss. A real deployment must measure each intermediate artifact because another ordering may work better for a different model or runtime.

Sequencing matters because these levers interact through the same weights and activations. Pruning changes the surviving distribution and tensor shapes, so subsequent quantization may

become easier or harder depending on ranges, outliers, formats, and kernels. Distillation can recover behavior in a chosen student, but it adds a separate objective and training run rather than guaranteeing recovery. Figure 10.30 reports one comparison from the cited compression study; it can motivate candidate sequences, but it does not establish a universal ordering among pruning, quantization, and singular value decomposition (SVD). Evaluate every intermediate artifact so a failure can be attributed to the transformation that introduced it.




Figure 10.30: Compression Methods in One Reported Setting: The cited comparison plots accuracy change against model-size ratio for pruning, quantization, their combination, and SVD. In this setting, pruning plus quantization reaches the smallest ratio near baseline accuracy; the ordering is model- and implementation-specific. Source: (Han, Mao, et al. 2016).

A mobile BERT pipeline makes that sequencing dependency concrete by combining pruning, distillation, and quantization under one deployment target.

10.8 Efficiency Measurement

Section 10.5 traced the gap between a compression ratio on paper and the speedup a model actually realizes. That gap now creates a measurement obligation. INT8 has up to $4\times$ smaller raw tensor payload than FP32, but latency must be measured on target hardware. Real speedups depend on memory hierarchy effects, kernel implementations, and hardware utilization patterns that theory alone cannot predict, so translating compression ratios into measurable improvements requires systematic profiling and evaluation. Three questions structure this analysis: *where* optimization efforts should focus, *how* to measure whether optimizations achieve their intended goals, and *how* to validate that combined techniques deliver expected benefits.

10.8.1 Profiling and opportunity analysis

Optimization begins with profiling to identify which components consume the most computational resources and offer the greatest optimization potential. The critical first step is determining whether model optimization will actually improve system performance, since model computation often represents only a fraction of total system overhead in production environments.

Modern machine learning models exhibit heterogeneous resource consumption: specific layers, operations, or data paths contribute disproportionately to memory usage, computational cost, or latency. Understanding these patterns is essential for prioritizing optimization efforts and achieving maximum impact with minimal accuracy degradation.

Effective profiling begins with establishing baseline measurements across relevant performance dimensions. Memory consumption, both static (model parameters and buffers) and dynamic allocation during inference, determines whether a model fits on the target device at all. Computational bottlenecks, measured in both FLOPs and actual wall-clock execution time, reveal which layers

dominate the inference budget. For battery-powered and edge deployments, power consumption profiles determine operational feasibility: a model that drains a phone battery in an hour is unusable regardless of its accuracy. End-to-end latency measurements identify which operations contribute most to inference delay, often revealing that memory-bound operations like layer normalization consume disproportionate wall-clock time relative to their FLOP count.

A critical caveat applies when translating profiling metrics into optimization estimates.

Consider a hypothetical Vision Transformer (ViT) profile for edge deployment. Suppose attention layers consume 65 percent of total FLOPs, layer normalization consumes 8 percent of latency despite only 2 percent of FLOPs, and the final classification head consumes 1 percent of computation but 15 percent of parameter memory. These scenario values motivate experiments rather than prescribe solutions. Candidate tests include hardware-supported structured pruning for attention, INT8 quantization of the classification head with accuracy validation, and runtime fusion around layer normalization.



Compression hits an end-to-end ceiling once non-model work dominates the pipeline.

Systems Perspective 10.5: FLOPs reduction is not proportional speedup

Reducing a model's FLOPs by 50 percent does not guarantee 50 percent latency reduction. Memory-bound operations see little benefit from arithmetic reduction because data movement, not multiplication, sets their time. Critically, Amdahl's Law applies at the system level (section D.2.3 derives its strong-scaling form): if model inference accounts for only 20 percent of end-to-end latency (with the remaining 80 percent spent on data loading, preprocessing, and postprocessing), then even perfect model optimization yields at most 1.25 $\times$ overall speedup. Tokenization and network transfer can occupy the non-model fraction in language-model serving; image decoding, resizing, and normalization can do the same in vision pipelines. Profile the complete path on target hardware before estimating optimization benefits.

Beyond these baseline measurements, modern optimization requires understanding model sensitivity to different types of modifications. Not all parameters contribute equally to accuracy. Layer-wise sensitivity analysis reveals which network components are most important for maintaining accuracy, guiding decisions about where to apply aggressive pruning or quantization and where to use conservative approaches.

10.8.2 Measuring optimization effectiveness


Optimization requires rigorous measurement frameworks that go beyond simple accuracy metrics to capture the full impact of optimization decisions. Effective measurement considers multiple objectives simultaneously: accuracy preservation, computational efficiency gains, memory reduction, latency improvement, and energy savings. Balancing these often-competing objectives requires careful trade-off analysis.

The measurement framework should establish clear baselines before applying any optimizations. Accuracy baselines must go beyond top-line classification accuracy to include calibration (whether confidence matches observed correctness), slice-level performance across important input or user groups, and robustness to input variations. Efficiency baselines capture computational cost (FLOPs, memory bandwidth), execution time across hardware platforms, peak memory consumption, and energy consumption profiles.

Quantizing ResNet-50 from FP32 to INT8 should be evaluated as a before-and-after system profile, not as an accuracy number alone. Table 10.16 uses scenario inputs, not measurements from one V100 benchmark, to illustrate the resource dimensions that determine deployment feasibility.

Because these values are illustrative, production decisions require the same measurements on the target workload and hardware.

With these comprehensive baselines in place, the measurement framework must track optimization impact systematically. Rather than evaluating techniques in isolation, applying our three-dimensional framework requires understanding how different approaches interact when combined. Sequential application can lead to compounding benefits or unexpected interactions that diminish overall effectiveness. Section 12.11.1.3 later expands this into a full efficiency-quality evaluation framework.



An illustrative INT8 profile must validate both size and accuracy.

Table 10.16: Illustrative ResNet-50 INT8 Deployment Profile: The values form a modeled scenario, not measurements from one hardware run. A real profile must measure accuracy, latency, artifact size, energy, and calibration on the target system.

Metric	FP32 baseline	INT8 result	Deployment reading
Top-1 accuracy	76.1%	75.8% (0.3 pp drop)	Aggregate accuracy mostly holds, but subgroup checks still matter.
Modeled V100 latency	1.56 ms	0.39 ms	GPU latency improves when the runtime maps INT8 to fast kernels.
Model size	102.4 MB	25.6 MB (4×)	The artifact becomes easier to cache, transmit, and deploy on edge.
Energy/inference	0.25 J	0.06 J (4×)	Lower precision reduces both arithmetic and memory-movement energy.
Calibration error	2.1%	3.4%	Confidence calibration can drift even when accuracy looks acceptable.

Rigorous measurement tells practitioners whether their optimizations succeeded, but the measurements themselves require tooling to perform. Profiling, quantization, pruning, and deployment all depend on software frameworks that automate otherwise prohibitively complex workflows.

10.9 Implementation Tools

Quantizing a 175-billion-parameter model by hand, inserting scale factors at every layer boundary, managing mixed-precision accumulation, and calibrating activation ranges would require modifying thousands of lines of model code. Without framework tooling, even straightforward INT8 post-training quantization demands manual insertion of quantization operations throughout the network, while pruning requires direct manipulation of weight tensors. Both become prohibitively complex as models scale.

Tool choice should follow the layer that owns the transformation. Framework APIs are useful while the model is still being trained, calibrated, or pruned because they can see weights, activations, gradients, and training state. Compiler and runtime tools become necessary after export, when the optimized graph must match a specific accelerator. This boundary is the practical reason software infrastructure matters: it records the calibration data, pruning schedule, quantization settings, and runtime libraries that produced the artifact, instead of leaving compression as an unrepeatable sequence of manual edits.

The resulting software infrastructure transforms theoretical optimization techniques into practical tools for production environments. For this chapter, the operational requirement is reproducibility. Chapter 14 later expands the same requirement into model versioning, monitoring, artifact management, and rollback procedures; here, the narrower point is that a compressed model must be traceable from training checkpoint to calibrated export to runtime artifact.

10.9.1 Model optimization APIs and tools

At the model-development layer, framework APIs are most useful for transformations that need access to training state, calibration data, or weight tensors before export. TensorFlow, PyTorch, and MXNet expose different APIs, but serve the same systems role. They insert quantization behavior into the graph, mutate or mask weights for pruning, and preserve enough metadata to reproduce the optimized artifact. Chapter 7 examines the frameworks themselves; this section only needs the compression-facing boundary.

Quantization-aware training is the clearest example because it must change the training graph without changing the mathematical objective. The transformation inserts quantization and dequantization behavior around selected tensors, lets gradients continue to flow through the simulated low-precision path, and records the configuration needed for later conversion. Listing 10.5 demonstrates the mechanism without depending on a specific framework API.

Listing 10.5: Quantization-Aware Training: Mechanism-level transformation that trains through simulated low-precision behavior and exports the calibration metadata needed for deployment.

```
choose target precision and calibration policy
insert quantize/dequantize boundaries around selected tensors
train with simulated low-precision values while keeping gradient flow intact
record scales, zero points, accumulator precision, and unsupported operations
export the calibrated graph and metadata as one reproducible artifact
```

The same transformation pattern applies to pruning. The optimization pass owns the weight tensor, applies a mask or structured removal rule, and records which parameters were removed so that subsequent fine-tuning and export operate on the intended model. Listing 10.6 illustrates the artifact boundary for both unstructured and structured pruning.

Listing 10.6: Pruning Artifact Boundary: Applies unstructured or structured removal rules, then records the mask and fine-tuning context needed to reproduce the compressed model.

```
choose pruning granularity: individual weights, channels, blocks, or attention heads
score candidate parameters by magnitude, sensitivity, or validation loss impact
remove or mask the selected structure according to the hardware-compatible rule
fine-tune the compressed model against the original validation target
export weights, masks, sparsity pattern, and calibration evidence together
```

The engineering value is repeatability rather than the API call itself. Built-in optimization APIs provide standardized control points for experimentation: teams can vary pruning amount, quantization mode, calibration data, and fine-tuning schedule while keeping the rest of the training pipeline fixed. That repeatability is what lets practitioners compare strategies rather than debug one-off graph rewrites.

10.9.2 Hardware-specific optimization libraries

After the model leaves the training framework, hardware-specific libraries own the final translation step: converting a pruned, quantized, or fused graph into kernels the target platform can execute efficiently. Libraries like TensorRT, XLA, OpenVINO, and TVM perform this translation for target platforms. Chapter 11 explains the accelerator features these tools target; the local point is that compression is incomplete until the exported graph maps to kernels the device can actually run well.

The runtime layer checks whether the apparent compression is executable compression. A pruned model needs sparsity-aware kernels or it may still run as dense arithmetic. A quantized model needs INT8 or INT4 kernels, scale handling, and layouts that the device supports. A fused model needs an operator pattern the compiler can legally combine without changing numerical behavior. These requirements explain why framework integration is not enough by itself: the artifact must survive conversion into a hardware-specific execution plan.

10.9.3 Diagnostic visualization

The toolchain boundary is therefore clear: framework APIs create and calibrate the compressed model, hardware optimization libraries adapt that model to the execution substrate, and diagnostic visualization asks whether the optimization damaged the model rather than only whether it shrank. The benchmarking and MLOps chapters later show how to validate and operate the exported artifact at system scale; within this chapter, the local diagnostic question is whether quantization, pruning, or sparsity changed the internal distributions in a way that threatens accuracy.

Quantization error histograms show whether error is broadly distributed or concentrated in outliers. Activation visualizations help detect clipping, overflow, and saturation. Figure 10.31 schematically groups first-layer convolutional kernels by visual pattern. Near-zero kernels or filters with consistently negligible activations—not merely uniform-looking filters—are pruning candidates, so the diagnostic must combine weight visualization with measured activation behavior. Framework

observers and debuggers can expose quantization ranges and tensor error, while compiler or runtime inspectors can show the graph and kernels that the exported model actually executes.




Figure 10.31: First-Layer Convolutional Filter Patterns: An 8×5 grid groups stylized kernels into edge, color, frequency, mixed, and uniform-looking patterns. Appearance alone does not establish inactivity; pruning requires measured activation or importance evidence in addition to low weight variance.

Sparsity diagnostics answer a deployment question: which layers actually became sparse enough for the runtime to exploit? Figure 10.32 illustrates how a layer-by-block heat map can reveal where sparsity concentrates, while trend plots track sparsity progression across pruning iterations. Kernel profiling must still determine whether the pattern yields speedup. TensorBoard, Netron, and SparseML provide these tools.

The diagnostic value is the nonuniformity, not the exact shade of any one cell. If later layers are much sparser than early feature extractors, the pruning policy is concentrating compression where representations are more redundant; if early layers darken first, the policy may be destroying low-level features before the model has enough depth to compensate.

Implementation tools make compression repeatable, but they do not make the trade-offs disappear. Quantization can follow pruning or distillation in a layered pipeline: pruning changes model structure or parameter count, quantization changes numerical representation, and distillation can train the smaller artifact against a teacher's behavior. These transformations can compound raw footprint reductions, but no fixed compression ratio or quality outcome transfers across models and deployments; each stage must be evaluated against the task target and profiled on the target execution path.

10.10 Fallacies and Pitfalls

The most instructive lessons in model compression often come not from what works but from what fails. The scenarios below illustrate common failure modes; exact outcomes depend on the model, compression method, workload, and target hardware.

Model optimization involves counterintuitive interactions between techniques that appear independent. Engineers often assume strategies compose linearly and that theoretical metrics predict deployment performance, and those assumptions waste optimization effort, degrade accuracy, or miss deployment requirements despite substantial investment.




Figure 10.32: Layer-Wise Sparsity Distribution: Block-level heat map visualizing non-uniform parameter removal. Darker tiles indicate higher sparsity density, revealing which network layers retain dense features versus which layers concentrate zero-valued weights for memory saving or sparse kernel scheduling.

Fallacy: *Optimization techniques can be applied independently without considering their interactions.*

Engineers assume optimization strategies compose additively: 50 percent pruning plus 4× quantization yields combined benefits. In reality, techniques interact nonlinearly and compound losses. In a representative BERT-style compression scenario, pruning to 70 percent sparsity may preserve most task performance, but applying INT8 quantization afterward can lose more accuracy than QAT on the pruned model. Knowledge distillation from heavily pruned teachers can also transfer degenerate attention patterns that reduce student accuracy compared with distilling from dense teachers. As section 10.7 demonstrates, successful optimization requires coordinated application where techniques are sequenced together. Organizations that apply aggressive combinations without measuring interactions waste weeks recovering lost accuracy.

Pitfall: *Optimizing for theoretical metrics rather than actual deployment performance.*

Teams can reduce FLOPs substantially without realizing proportional latency gains. In one hypothetical trace, a pruned model with 40 percent fewer parameters achieves only 12 percent latency reduction because irregular sparsity prevents efficient execution. A second scenario reduces a transformer from 440 MB to 110 MB, yet conversion overhead on hardware without a fast low-precision path erodes the latency benefit. These values illustrate failure modes rather than report ARM or GPU measurements. Memory bandwidth, cache behavior, kernels, and instruction-level parallelism determine actual performance, so production deployments require wall-clock measurements on target hardware.

Fallacy: *Aggressive quantization maintains model performance with minimal accuracy loss.*

Engineers assume quantization error scales uniformly with bit width. In practice, precision reduction can exhibit threshold effects, and the threshold depends on the model, layer, quantizer, calibration,

and training method. The ResNet and BERT values in this chapter are illustrative rather than universal degradation bands. Even operations such as LayerNorm and Softmax can be implemented with integer approximations, so they do not universally require FP16; whether a low-precision implementation preserves task quality must be measured. As section 10.4.4 demonstrates, mixed-precision approaches can retain higher precision where uniform quantization fails.

Pitfall: *Defaulting to FP32 everywhere to avoid quantization risk.*

FP32 uses twice the raw storage and bandwidth of BF16, but lower-precision execution does not automatically preserve convergence or accuracy. Mixed-precision training often retains FP32 accumulators or master weights and may require loss scaling, while inference precision must be validated per model and operator. The right question is what lowest precision preserves quality on the evaluation distribution and maps efficiently to the target hardware.

Fallacy: *Post-training optimization is always enough.*

Teams often begin with post-training quantization (PTQ) because it avoids retraining. If the resulting artifact meets its task and deployment thresholds, that simple path may be sufficient. When it does not, quantization-aware training (QAT) can recover part of the quality gap by adapting the model to simulated low-precision behavior. Posttraining pruning and post-hoc distillation can likewise underperform training-aware schedules at aggressive sparsity or bit width. As detailed in section 10.4.4, the production threshold—not a presumption that either path is universally superior—governs the decision.

Pitfall: *Assuming compression ratios translate directly into proportional deployment gains.*

Teams may obtain a 4× raw weight-payload reduction through INT8 quantization and expect the same deployment gain. In practice, scales, metadata, unsupported operators, conversions, and sparse indexing erode the benefit. The chapter’s 18.8 percent end-to-end improvement versus an 8× paper target is a hypothetical trace, not a BERT benchmark. Production workflows must profile deployed latency on target hardware rather than extrapolate from compression ratios.

Fallacy: *Sparse matrices always save memory.*

Sparse formats add index-array metadata that erodes savings at moderate densities. With FP32 values and INT32 indices, CSR approaches break-even near 50 percent density before row-pointer overhead, while COO stores additional coordinates and crosses over at a lower density. Performance has no universal sparsity threshold because tensor shape, format, kernel, batch, and hardware determine whether sparse execution outruns dense execution. Specialized formats such as NVIDIA’s 2:4 sparsity reduce metadata and provide a supported execution path, while arbitrary unstructured sparsity may deliver neither memory nor latency savings on commodity hardware (see section 10.8.1).

Pitfall: *Choosing sparse storage before checking the density threshold.*

Teams sometimes convert tensors to sparse formats as soon as pruning creates visible zeros. That conversion can make the system slower and larger if metadata, gather/scatter overhead, and poor cache locality outweigh the saved values. A production compression pass should measure the realized density, choose a sparse format only after the break-even point is crossed, and prefer structured sparsity when the target hardware has kernels that can exploit it.

10.11 Summary

Model compression is not a bag of tricks but an engineering discipline built on three complementary dimensions: *structural optimization* determines what the model computes, *precision optimization* determines how precisely it computes, and *architectural optimization* determines how efficiently those computations execute on physical hardware. These dimensions can compound, but their gains do not multiply automatically. Pruning, distillation, and quantization together produce the illustrative pipeline’s 16× footprint reduction from 440 MB to 28 MB; metadata, unsupported operators, and hardware mapping determine the realized deployment gain. A parameter reduction approaches the same latency reduction only when the affected work is on the critical path and the runtime can exploit the new representation. Profile on target hardware, not paper metrics.

Combined with the data selection techniques from chapter 9, these model-centric optimizations complete the model-side efficiency toolkit: data selection maximizes learning from available exam-

ples, while model compression minimizes resources required for deployment. Whether those savings become real speedups still depends on the hardware and runtime that execute the compressed model.



Key Takeaways: From benchmark winner to production model

- **Compression spends surplus capacity:** Production models trade unused parameters, precision, and capacity for latency, memory, power, or cost limits the deployment cannot violate. The right target is not minimum size, but the smallest artifact that preserves the behavior the context actually needs.
- **Savings multiply only when aligned:** Structural pruning, distillation, quantization, and architecture changes can compound raw footprint reductions, as the modeled BERT mobile pipeline's $16\times$ ratio illustrates. The gain becomes real only when task quality holds and the resulting operators match the target runtime and accelerator.
- **Precision is a deployment contract:** Where the target runtime supports it, INT8 post-training quantization is a useful first experiment because it reduces raw FP32 weight payload by $4\times$ without retraining. QAT, distillation, or mixed precision become candidates when calibration or layer sensitivity exposes unacceptable error.
- **Hardware sets the exchange rate:** Unstructured sparsity and theoretical FLOP cuts do not imply latency gains unless kernels and memory layouts can exploit them. The chapter's hypothetical $8\times$ paper target versus $1.5\times$ realized outcome illustrates why target-hardware profiling is mandatory.
- **End-to-end latency caps model wins:** Compression is valuable only on the critical path. When inference is 20 percent of total request latency, Amdahl's Law caps even perfect model acceleration at $1.25\times$, so preprocessing, dispatch, and data movement may be the true optimization target.

The techniques in this chapter differ in almost every detail, yet one rule runs beneath all of them. Each buys a smaller, faster, or cooler model by spending something else: pruning spends capacity, quantization spends numerical precision, and distillation spends training compute. Where the model holds no surplus to spend, the bill may be paid in task quality. Compression does not make the work disappear; it moves cost or risk from one part of the system to another, and the hardware sets the exchange rate. A technique that pays off on a phone can therefore be worthless on a data-center GPU. This is the conservation-of-complexity heuristic at the level of a single model. A production artifact is not merely a benchmark model made smaller; it is a model whose structure, representation, and execution path have been engineered together until the trade balances against the silicon and application constraints that govern deployment.




What's Next: From math to physics

These optimization techniques compress the model's logic, reducing the work the algorithm asks the machine to perform. Logic, however, must eventually run on physical silicon. Chapter 11 examines how GPUs, TPUs, and NPUs—with their systolic arrays, Sparse Tensor Cores, and tiered SRAM/HBM memory hierarchies—map supported compressed representations onto the throughput and data movement available in hardware.

11

Hardware Acceleration




- 11.1 Acceleration Fundamentals
- 11.2 Hardware Specialization
- 11.3 AI Compute Primitives
- 11.4 Compute Units and Execution Models
- 11.5 AI Memory Systems
- 11.6 Roofline Model
- 11.7 Hardware Mapping
- 11.8 Dataflow Optimization
- 11.9 Compiler Support
- 11.10 Runtime Support
- 11.11 Multi-Chip Scaling
- 11.12 Heterogeneous SoC Design
- 11.13 Hardware Sustainability
- 11.14 Fallacies and Pitfalls
- 11.15 Summary

Purpose

Why does moving data cost more than computing it?

Arithmetic is inexpensive relative to many memory accesses. In the time it takes to fetch a value from main memory, a processor can perform many calculations. This imbalance, the “memory wall,” reflects the latency, bandwidth, and energy costs of moving data across a system. It explains why specialized accelerators are not merely faster at math but are architected to hide, amortize, and reduce the cost of moving data through deep memory hierarchies, massive parallelism, and specialized data paths. Hardware acceleration allows many large AI workloads to grow when general-purpose processor scaling alone is insufficient. It also explains why some optimizations that reduce theoretical computation fail to improve runtime. If an operation is memory bound, computing less may not help because computation is not the bottleneck. Hardware selection therefore cannot be reduced to comparing peak FLOP/s. What matters is whether a workload’s data movement patterns align with what the hardware was designed to accelerate. A model with large embedding tables and irregular lookups needs a different accelerator than one performing dense matrix multiplications over compact weight tensors. Software must expose and schedule the available reuse; otherwise specialized units wait while data moves and headline throughput remains theoretical. Matching workload and hardware determines whether execution remains at a fraction of theoretical peak or approaches the hardware’s practical ceiling. In D·A·M terms, the accelerator makes the machine constraint concrete and shapes which algorithms remain practical in production.






Learning Objectives

- Explain hardware acceleration as machine-axis specialization for tensor workloads, data reuse, and performance per watt
- Calculate arithmetic intensity and roofline ceilings to classify kernels as compute bound or memory bound
- Diagnose memory-wall bottlenecks using bandwidth, cache hierarchy, host-device transfer, and energy-movement costs
- Compare Tensor Cores, systolic arrays, SIMD/SIMT units, and sparse execution for ML compute primitives
- Select dataflow, tiling, and mapping strategies that maximize reuse under memory-capacity constraints
- Analyze compiler and runtime optimizations that fuse kernels, plan memory, and schedule accelerators
- Evaluate accelerator choices across throughput, latency, power, cost, and deployment-context constraints

11.1 Acceleration Fundamentals



Hardware acceleration turns on the machine axis.

REDUCING parameters, precision, or operations only matters when the machine can execute the resulting representation efficiently. Data selection reduced the data term, and compression reduced the algorithm's work; hardware acceleration asks what the machine can actually deliver. The answer starts with the memory wall: arithmetic is often cheap relative to data movement. An off-chip memory access can take many processor cycles and consume orders of magnitude more energy than a low-precision arithmetic operation, with the exact ratio depending on the device and memory level. Specialized hardware matters because it raises compute throughput while organizing memory, dataflow, and parallelism so those arithmetic units stay fed.



Definition 11.1: Hardware acceleration

Hardware acceleration is the practice of mapping work from general-purpose execution onto specialized processors or functional units optimized for a narrower operation class. For ML workloads, that specialization commonly trades some programmability for the compute density (R_{peak}) and performance-per-watt gains that regular tensor operations can exploit.

1. **Significance:** An A100 GPU delivers 312 TFLOP/s for FP16/BF16 matrix multiplication; against this chapter's illustrative 1–2 TFLOP/s reference-CPU baseline, this is a 156–312× gap. Its 54.2 billion transistors devote far more silicon to parallel arithmetic than to branch prediction, out-of-order scheduling, and large caches (NVIDIA Corporation 2020a; Choquette et al. 2021).
2. **Distinction:** A general-purpose CPU is designed for flexible execution, low-latency control, and varied workloads. An accelerator instead concentrates resources on a narrower operation class, so it achieves its gains only when the workload exposes the supported operations and enough parallel work to keep its arithmetic units busy.
3. **Common pitfall:** A frequent misconception is that an accelerator's advertised peak throughput is the throughput a workload receives. Delivered performance is the lower of the compute ceiling and what memory bandwidth can feed, the roofline constraint: a low-arithmetic-intensity kernel can sit at a small fraction of peak FLOP/s no matter how fast the silicon is rated, because it starves for data rather than for arithmetic.

This definition frames the chapter's central engineering trade-off. General-purpose processors support broad control flow, memory, and execution patterns. An accelerator can devote more of its design budget to the arithmetic, storage, and data paths of a narrower workload class. The result

can be order-of-magnitude improvements in throughput per watt when the workload and software path match that specialization.

Hardware alone, however, cannot achieve these gains. The algorithms must be designed to exploit what the hardware offers, and the hardware must be built to accelerate the operations algorithms actually use. This symbiosis motivates a complementary principle: hardware-software co-design.

Definition 11.2: Hardware-software co-design

Hardware-Software Co-design is an ML accelerator development method that crosses conventional abstraction boundaries, allowing algorithm constraints to inform silicon design and hardware capabilities to shape algorithm formulation.

1. **Significance:** On a supported tensor-core path, 8-bit integer (INT8) quantization can deliver multi-fold throughput gains because hardware executes lower-precision operations more densely and moves fewer bytes than an FP32 path; the exact gain depends on the architecture and workload (NVIDIA Corporation 2020a; Dally et al. 2021; Dally 2023).
2. **Distinction:** Unlike layered abstraction (where software calls a hardware API without knowing the silicon details), co-design exposes hardware constraints directly to algorithm and compiler authors: data alignment requirements, precision formats, and memory access patterns all become visible inputs to global cross-layer optimization.
3. **Common pitfall:** A frequent misconception is that co-design is a one-time hardware design choice. In practice, the supported features continue to evolve: NVIDIA Tensor Cores began with FP16 matrix multiplication, later generations added TF32 and broader integer support, and Ampere added acceleration for 2:4 structured sparsity (NVIDIA 2017; NVIDIA Corporation 2020a).

Co-design explains *why* the compression techniques introduced in chapter 10 deliver real speedups. The quantization techniques in section 10.4 show why converting FP32 to INT8 can accelerate a supported hardware path: fewer bytes reduce traffic, and the accelerator may execute more low-precision operations per cycle than FP32 operations (NVIDIA Corporation 2020a). Structured pruning can improve performance when it produces smaller dense shapes or a sparsity pattern supported by the hardware, whereas arbitrary unstructured sparsity needs matching sparse kernels and metadata handling. Evaluating accelerator performance requires tracing the path from workload to silicon: compute primitives, memory systems, roofline diagnosis, mapping and dataflow, then compiler and runtime support. The recurring question is why some promising algorithmic optimizations survive contact with hardware while others remain paper savings.

Theorem 11.1: The fundamental limit of acceleration (Amdahl's Law)

Hardware acceleration *does not speed up the entire system*; it speeds only the baseline-time fraction using the accelerated path. For that fraction p and path gain G_{accel} , **Amdahl's law for AI** (Amdahl 1967) gives the bound in equation 11.1.

$$\text{Speedup} = \frac{1}{(1-p) + \frac{p}{G_{\text{accel}}}} \quad (11.1)$$

- **Accelerated fraction (p):** Share of baseline time receiving the stated gain, often in matrix multiplication or another supported kernel.
- **Accelerator gain (G_{accel}):** The raw speed advantage of the GPU or Tensor Processing Unit (TPU) over the CPU for the accelerated portion of the workload.
- **Unaccelerated fraction ($1-p$):** Work outside the chosen path, such as data loading, host overhead, unsupported operators, and kernel launch latency.

Pitfall: *Unaccelerated work caps total speedup.* If 10 percent of baseline time is unaffected ($p = 0.9$), even infinite path gain ($G_{\text{accel}} = \infty$) yields at most $10\times$ overall. Once the accelerated component is fast enough, unaffected work dominates.

Hardware acceleration targets specific terms in the iron law of ML systems (section 1.7), which decomposes end-to-end time into data volume (D_{vol}/BW), computation ($O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$), and fixed latency (L_{lat}). While data selection reduced the total data and model compression reduced O per sample, hardware acceleration increases the rate at which those operations execute by improving R_{peak} , η_{hw} , and BW. Section D.2 supplies the analytical performance models that diagnose which of these terms dominates a given workload, including the dimensional analysis that confirms each iron law term resolves to seconds. Yet acceleration has a hard ceiling, established by Amdahl's Law.¹

Amdahl's Law also explains *why* many GPU upgrades disappoint. Figure 11.1 visualizes the **acceleration wall**: when meaningful baseline time remains unaffected, the ceiling $1/(1-p)$ makes further gain progressively less valuable unless p approaches one. The contours are illustrative.

¹ | **Amdahl's law:** Maps onto the iron law only after the proposed change is specified. If a faster tensor unit reduces the computation term ($O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$) but leaves data transfer and fixed latency unchanged, those unaffected terms bound total speedup. A different accelerator or memory system may improve bandwidth as well. The unaccelerated fraction is therefore defined by the intervention, not by a permanent classification of each term as serial.




Figure 11.1: The Iron Law Heatmap: Total system speedup as a function of accelerator gain (G_{accel}) and accelerated fraction (p). High speedup appears only near the top-right corner, where both accelerator gain and accelerated fraction are high. At $p = 0.9$, for example, even an infinite gain on the accelerated path cannot exceed $10\times$ end-to-end speedup. Contours span roughly $1\times$ – $500\times$ speedup.

The key intuition to carry into specific hardware architectures is that raw speedups matter only after the unaccelerated fraction has been reduced. The accelerated fraction p can differ substantially between workload archetypes running on the same hardware, and at fleet scale these differences determine whether an accelerator investment pays off or stalls at the remaining bottleneck.



Checkpoint 11.1: The parallelism gate

These checks turn Amdahl's bound into a hardware-selection test.

Amdahl's reality

- ☐ **Serial bottlenecks:** use Amdahl's bound to explain why a $1,000\times$ faster GPU may only speed up training by $5\times$ when data loading is slow.
- ☐ **Speedup ceiling:** for $p = 0.95$ and $G_{\text{accel}} = 100$, estimate total speedup and compare it with the infinite-accelerator ceiling.
- ☐ **Workload variation:** compare ResNet-50 and MobileNet accelerated fractions for a proposed hardware change, then predict which benefits more.

The serial bottleneck becomes concrete on real hardware, where the same accelerator that nears its parallel ceiling on one workload can stall on another. Section D.1 collects the reference R_{peak} figures across accelerator generations and the latency hierarchy that ground these hardware comparisons in order-of-magnitude terms.

Lighthouse 11.1: An illustrative Amdahl budget on H100

Illustrative ResNet-50 inference on NVIDIA H100:

- Assume H100 provides $G_{\text{accel}} = 247\times$ on matrix multiplication relative to a reference CPU without matrix extensions (1979 TOPS INT8 vs. ~ 8 TOPS) (Choquette 2023).
- Assume $p = 0.95$: 95 percent of baseline time receives that gain, while 5 percent remains in data loading, preprocessing, postprocessing, and other work.

$$\text{Speedup} = \frac{1}{(1 - 0.95) + \frac{0.95}{247}} = \frac{1}{0.05 + 0.0038} \approx 18.6\times$$

The $247\times$ path gain produces only $18.6\times$ end-to-end speedup because the 5 percent unaffected fraction sets the ceiling.

Contrast with GPT-2 (autoregressive):

- For GPT-2 token generation, assume $p = 0.80$; the remaining 20 percent includes host overhead, sampling, and decoding work outside the accelerated matrix path.

$$\text{Speedup} = \frac{1}{(1 - 0.80) + \frac{0.80}{247}} = \frac{1}{0.20 + 0.0032} \approx 4.9\times$$

Less token-generation time receives the assumed gain, so even an infinite path gain yields at most $1/(1 - p) = 5\times$. Serving optimizations address this limit differently: batching improves cross-request utilization, while speculative decoding uses a draft model to reduce sequential target-model steps.

These examples reveal that hardware optimization turns on whether a workload is limited by compute rate or data movement. That distinction determines which accelerator to choose, which optimizations matter, and whether a $10\times$ more powerful chip will help. The Roofline Model provides the analytical framework for making this diagnosis (Williams et al. 2009); it is introduced formally in section D.2.1 and section 11.6 applies it to AI workloads. It plots an operation's arithmetic intensity,² defined as the ratio of floating-point operations to bytes of memory traffic (FLOP/byte), against hardware capabilities, revealing whether performance is capped by compute or bandwidth. A dense matrix multiplication with high arithmetic intensity benefits from more TFLOP/s; a LayerNorm with low arithmetic intensity benefits from more memory bandwidth. High-reuse ResNet-50 convolutions can cross into the compute-bound regime, while low-batch autoregressive attention can be memory bound. This distinction is precisely *why* these workloads require different optimization strategies.

The chapter develops these ideas from hardware history through computational primitives, memory hierarchies, roofline analysis, mapping, dataflow, and runtime support. The core analysis stays with single-accelerator and single-node systems; the closing material uses multi-device examples only to show how the same bottleneck diagnoses scale. The history of specialized hardware comes first because it reveals the recurring design patterns behind modern accelerators.

11.2 Hardware Specialization

The TPUv1/K80 efficiency shock is a modern instance of a recurring hardware pattern: when a workload becomes important and regular enough, specialized hardware can outperform a more general platform. Machine learning acceleration follows a trajectory also seen in floating-point arithmetic, graphics processing, and digital signal processing. Each era specialized the operation or data path that constrained its workload; in modern ML, the cost of feeding dense arithmetic makes data movement a central design target.

² | **Arithmetic intensity:** The ratio of compute operations to bytes transferred across the memory interface being modeled (FLOP/byte). Workloads above that hardware level's ridge point, such as well-tiled matrix multiplications with substantial reuse, are compute bound and can benefit from more TFLOP/s. Low-intensity operations, including some low-batch attention and decoding kernels, are bandwidth bound at that level and need more bandwidth or reuse rather than only more compute throughput.

3 | TPU (tensor processing unit):

The first TPU made a deliberately narrow bet, centering the design on a single 256×256 systolic array for 8-bit matrix multiplication and avoiding much of the control machinery required by a general-purpose core (Jouppi et al. 2017). That trade buys high compute density on dense matrix multiplication at the cost of flexibility. Irregular or branch-heavy code remains a poor fit, and compiler tiling, padding, and layer dimensions determine how fully the array is used.

Modern ML accelerators (DianNao-class neural-network accelerators (Chen et al. 2014), GPUs with Tensor Cores, Google’s TPUs,³ Apple’s Neural Engine) emerged from these established architectural principles. The evolution spans four phases: specialized computing origins, parallel graphics processing, domain-specific architectures, and the emergence of ML-specific hardware. Each phase reveals design principles that remain relevant for understanding and optimizing contemporary AI systems.



Example 11.1: The TPUv1 vs. K80 efficiency shock

Scenario: In 2015, Google deployed its first Tensor Processing Unit (TPUv1) into production data centers, comparing its throughput and energy efficiency against general-purpose NVIDIA K80 GPUs (Jouppi et al. 2017).

Diagnosis: The K80 was a programmable GPU designed to support a broader mix of graphics and compute workloads. The TPUv1 domain-specific architecture (DSA) made a narrower inference-focused bet, pairing a large 8-bit systolic matrix unit with software-managed memory and simpler control.

Systems lesson: On the six Google inference workloads and contemporary CPU/GPU baselines evaluated in the TPU paper, tailoring silicon to the dominant matrix operations delivered $15\text{--}30\times$ higher throughput and $30\text{--}80\times$ better performance per watt (Jouppi et al. 2017). Those ratios describe the evaluated systems, not every workload.

Hardware specialization improves performance by implementing frequent patterns in dedicated circuits, but introduces trade-offs in flexibility, silicon area, and programming complexity. The principles that shaped early floating-point and graphics accelerators now inform AI hardware design.

11.2.1 Specialized computing

Hardware specialization emerges when specific computational patterns become the primary system bottleneck, preventing general-purpose processors from scaling efficiently. This history highlights three recurring limits: slow implementation of scalar floating-point, insufficient throughput for parallel graphics, and poor integration between memory and parallel compute.

The first phase, the precision bottleneck, occurred when scientific and engineering applications required floating-point arithmetic that general-purpose CPUs performed poorly. In the late 1970s, CPUs typically emulated floating-point operations in software, requiring many instructions for a single multiplication. This scalar inefficiency led to the first major instance of hardware specialization: the mathematics coprocessor.

The Intel 8087 (1980)⁴ addressed this bottleneck by offloading arithmetic-intensive tasks to a dedicated unit. Implementing floating-point logic in hardware greatly accelerated operations that otherwise required software emulation (Palmer 1980). This established a core principle: moving a dominant operation to specialized silicon can provide a substantial speedup.

As specialized functions like floating-point math proved their value, they followed a recurring pattern of integration. The Intel 486DX (1989) moved the FPU directly onto the CPU die, eliminating the off-chip communication latency and making high-precision math a standard feature rather than an optional accelerator (Hennessy and Patterson 2017). This cycle, specialization to solve a bottleneck followed by integration into the general-purpose stack, has recurred across several eras of hardware evolution.

The progression from specialization to integration has shaped modern computing. Each domain (graphics, signal processing, machine learning) introduced specialized architectures that were later absorbed into general-purpose platforms.

Figure 11.2 traces this recurring cycle of specialization and integration across five eras, each addressing the dominant computational bottleneck of its period: the 1980s floating-point and signal-processing units (Intel 8087, TI TMS32010 DSP), 1990s 3D graphics (NVIDIA GeForce 256), 2000s media and network processing (H.264 codecs, Intel IXP2800), 2010s deep-learning tensor operations (Google TPU v1, NVIDIA Tensor Cores), and 2020s application-specific accelerators (AI engines,

4 | **Intel 8087:** The coprocessor implemented floating-point logic directly in silicon, avoiding the CPU’s slow, multi-instruction software emulation for each calculation. Its benefit depended on how much time an application spent in supported floating-point operations, illustrating why specialization helps most when the accelerated work dominates execution.

wafer-scale ML chips). Capabilities such as real-time translation, recommendations, and on-device inference build directly on principles established in these earlier specialization waves.




Figure 11.2: Hardware Specialization Timeline: Computing architectures evolve through recurring cycles of specialization and integration across five eras. Each era introduces domain-specific hardware to resolve an important memory or arithmetic bottleneck—from off-chip FPUs to GPUs, codecs, and TPU/Tensor Core array units. Many of these capabilities later appear alongside CPUs in broader systems, even when the specialized units remain distinct.

11.2.2 Parallel computing and graphics processing

The principles established through floating-point acceleration provided a blueprint for addressing subsequent computational challenges. As computing applications diversified, new computational patterns emerged that exceeded the capabilities of general-purpose processors, and each domain contributed unique insights to hardware acceleration strategies.

Graphics processing emerged as a primary driver of hardware specialization in the 1990s. Early graphics accelerators focused on specific operations like bitmap transfers and polygon filling. NVIDIA's GeForce 256 in 1999 represented a milestone in specialized computing. The GeForce 256 implemented hardware-accelerated transform and lighting (T&L), moving these computations from CPU to dedicated silicon. While not yet programmable, these Graphics Processing Units (GPUs) demonstrated how fixed-function parallel architectures could efficiently handle data-parallel workloads such as texture mapping and vertex transformation. The transition to programmable shaders with the GeForce 3 (2001) and unified shader architectures with the GeForce 8 (2006) eventually enabled GPU computing for general-purpose workloads. By 2004, high-end GPUs could process over 100 million polygons per second (Owens et al. 2008).

Concurrently, Digital Signal Processing (DSP) processors established parallel data path architectures with specialized multiply-accumulate units and circular buffers optimized for filtering and transform operations. Texas Instruments' TMS32010 (1983) demonstrated how domain-specific instruction sets could dramatically improve performance for signal processing applications (Lyons 2011).

Network processing introduced additional patterns of specialization. Network processors developed unique architectures to handle packet processing at line rate, incorporating multiple processing cores, specialized packet manipulation units, and tiered memory management systems. Intel's IXP2800 network processor shows the consequence of one hard constraint: meeting line-rate packet deadlines leaves no slack for cache misses, so the design arranges many parallel cores around tiered on-chip memory to keep data adjacent to compute. That compute-near-memory organization, forced here by packet timing, is the same arrangement ML accelerators later adopt to keep their processing-element grids fed.

Across these domains, a common blueprint emerges: identify the dominant computational patterns, build specialized processing elements and memory hierarchies around them, create tailored programming models, and progressively evolve toward more flexible architectures. This pattern of architectural co-evolution established the foundation for contemporary AI hardware design. DSP

5 | AlexNet: Krizhevsky, Sutskever, and Hinton's 60-million-parameter convolutional neural network (CNN) won ImageNet 2012 by more than ten percentage points on two consumer GTX 580 GPUs with only 3 GB of VRAM each. Because the model exceeded single-GPU memory, the authors partitioned the network across two cards and limited communication between them, an early form of model parallelism that foreshadowed systematic tensor and pipeline parallelism. The reported training run took five to six days (Krizhevsky et al. 2012).

6 | DSA (domain-specific architecture): Hardware optimized for an application domain, sacrificing some general-purpose programmability for efficiency. Google's TPUv1 achieved 15–30 $\times$ better performance and 30–80 $\times$ better performance per watt than the contemporary CPU and GPU systems evaluated on Google's inference workloads by centering its design on a systolic array and software-managed memory (Jouppi et al. 2017). A DSA that excels at dense matrix multiplication may perform worse than a CPU on irregular workloads such as graph traversal, so the workload limits and efficiency gain must justify the software, compiler, and ecosystem cost (Hennessy and Patterson 2019, 2017).

7 | Moore's law: In the illustrative assumptions used by figure 11.3, model compute demand grows roughly 6.1 $\times$ per year while accelerator peak supply improves roughly 1.7 $\times$ per year, widening the modeled demand/supply gap by about 3.5 $\times$ per year. Published compute trends motivate the divergence, but the plotted rates are assumptions rather than forecasts. The figure therefore shows how sensitive the systems gap is to sustained differences in growth rates, not a fixed future trajectory. Under such a gap, algorithmic efficiency techniques such as compression, quantization, and sparsity become increasingly valuable (Amodei and Hernandez 2018a; Epoch AI 2024).

innovations in low-power signal processing enabled real-time inference on edge devices, including voice assistants and wearables. Together, these domains informed ML hardware designs and demonstrated that accelerators could be deployed across both cloud and embedded contexts.

A single result made the GPU's relevance to AI unmistakable. AlexNet⁵ (Krizhevsky et al. 2012) won the 2012 ImageNet competition by more than ten percentage points on two consumer-grade NVIDIA GTX 580 graphics cards, each with only 3 GB of VRAM. The systems lesson was that matching a workload's data parallelism to GPU hardware could make previously impractical training runs feasible. The era of GPU-centric deep learning had begun.

11.2.3 Emergence of domain-specific architectures

These diverse acceleration patterns converged in a broader architectural shift. The emergence of **domain-specific architectures (DSAs)**⁶ marks a transition in computer system design, driven by the slowing benefits of traditional scaling (Esmailzadeh et al. 2011) and the increasing computational demands of specialized workloads. Moore's Law⁷ described a long-running increase in transistor density (Moore 1998). Dennard scaling⁸ (Dennard et al. 1974) allowed voltage to fall as transistors shrank; its breakdown removed the accompanying path to routine frequency gains at constant power density. Together, these shifts constrained general-purpose performance and efficiency. As Hennessy and Patterson (2019) argued in the 2017 Turing Lecture, the response was a new era of domain-specific solutions optimized for important workloads.

The scale of this challenge becomes stark in figure 11.3, which plots an illustrative systems gap between model demand and hardware supply, sometimes discussed under the informal label Huang's Law.⁹ The normalized scenario assumes GPU supply rises about 1.7 $\times$ per year while model demand rises about 6 $\times$ per year, producing a gap that widens by roughly 3–4 $\times$ each year. Published compute trends motivate the qualitative divergence, while the plotted rates are scenario assumptions rather than a forecast (Amodei and Hernandez 2018a; Epoch AI 2024).




Figure 11.3: The Systems Gap: Relative compute growth (log scale) comparing model demand to hardware supply, normalized to 2012 = 1.0. The gray dotted line (CPU) and blue dashed line (GPU) reflect hardware progress, which lags the exponential red solid line (Model Demand). The purple region is the 'Systems Gap' that must be bridged through parallelism and co-design.

The plot is normalized to a 2012 baseline to emphasize relative growth. Notice how the purple-shaded region between the curves keeps widening—this gap cannot be closed by waiting for faster chips; it requires architectural innovation.

11.2.4 The technology S-curve: Why we must shift

Many computing technologies are described by a lifecycle of ferment (initial slow progress), take-off (rapid growth), and saturation (diminishing returns near physical or economic limits). Figure 11.4 uses two conceptual S-curves to express the pattern: as gains from one design approach taper, domain-specific architectures can open a new efficiency curve for workloads with stable computational structure.




Figure 11.4: The Twin S-Curves of Specialized Computing: This conceptual diagram represents a general-purpose design approach (gray) reaching diminishing returns as a specialized approach (blue) begins a new efficiency curve. It is a schematic of an architectural transition, not a measured historical performance series. Specializing hardware for stable linear-algebra workloads can produce large gains at the cost of general programmability.

The routine gains once associated with shrinking transistors have slowed, while some measures of AI model demand have grown much faster than hardware supply. The response cannot be limited to waiting for the next CPU generation; it requires architectural changes that improve how important workloads use transistors, memory bandwidth, and energy. Understanding this transition requires examining the scaling laws that fueled the general-purpose era.

Historically, improvements in processor performance depended on semiconductor process scaling and increasing clock speeds. As power density limitations restricted further frequency scaling and transistor miniaturization encountered increasing physical and economic constraints, architects explored alternative approaches to sustain computational growth. The result was a shift toward domain-specific architectures, which dedicate silicon resources to optimize computation for specific application domains, trading flexibility for efficiency.

Domain-specific architectures achieve superior performance and energy efficiency when the hardware stops treating the workload as arbitrary code. The first shift is a customized data path: matrix multiplication units in AI accelerators, for example, implement **systolic arrays**, grid-like networks of processing elements that rhythmically compute and pass data through neighboring units. Once that data path is fixed, the memory hierarchy can be tuned around the reuse pattern the workload actually needs, with cache configurations, prefetching logic, and memory controllers designed for the expected tensor flow.

The same specialization then reduces control overhead. Domain-specific instruction sets encode common operation sequences into single instructions, minimizing decode and dispatch complexity, while fixed-function circuit blocks bypass software interpretation for operations that appear constantly. The result is not one trick but a stack of matching decisions: data movement, memory locality, instruction overhead, and circuit implementation all align around the same computational pattern.

Modern smartphones illustrate these principles compellingly. They can decode high-resolution video within tight power and thermal envelopes even though video processing requires billions of operations per second. This efficiency is achieved through dedicated hardware video codecs¹⁰ that implement industry standards such as H.264/AVC and H.265/HEVC (Sullivan et al. 2012). These specialized circuits can provide order-of-magnitude performance-per-watt gains compared with software decoding on general-purpose processors, with the exact gain depending on codec, resolution, process node, and CPU baseline.

These later domains are not separate anecdotes; they repeat the same bottleneck response. Genomics processing benefits from custom accelerators because long-read assembly contains stable alignment kernels that dedicated hardware can accelerate (Turakhia et al. 2018). Blockchain computation produced application-specific integrated circuits (ASICs)¹¹ for the same reason: cryptographic hashing is fixed enough to justify silicon that trades flexibility for efficiency (Bedford Taylor 2017).

⁸ | **Dennard scaling:** The 1974 scaling model held that operating voltage could fall as transistor dimensions shrank, keeping power density approximately constant (Dennard et al. 1974). Its breakdown after about 2005 constrained further clock-frequency growth within a chip's thermal design power; the resulting **dark silicon** literature estimated, for modeled advanced nodes and design assumptions, that thermal limits could prevent large fractions of available transistors from operating simultaneously. That fraction is not a universal constant, but the power constraint is durable: specialized units make the active silicon budget more useful by dedicating it to high-value operation classes (Esmaeilzadeh et al. 2011).

⁹ | **Huang's law:** The observation that GPU performance for AI workloads has improved through architectural innovations such as Tensor Cores as well as transistor scaling. The normalized figure uses illustrative GPU-supply and model-demand curves to show how unequal growth rates create a widening systems gap; it is not a forecast of either rate (Amodei and Hernandez 2018a; Epoch AI 2024; NVIDIA Corporation 2020a; Choquette 2023).

10 | Codec: A portmanteau of “coder-decoder,” reflecting the hardware’s dual function. Encoding (compression) is compute-intensive because it searches for optimal representations, while decoding (decompression) is bandwidth-intensive because it reconstructs full-resolution frames from compressed streams. Dedicated codec silicon implements both paths in fixed-function hardware, so neither path wastes transistors on unrelated general-purpose control logic.

11 | ASIC (application-specific integrated circuit): These circuits achieve their extreme efficiency by implementing a single algorithm directly in silicon, often improving performance-per-watt by $10^3\times$ to $10^5\times$. Examples include cryptographic hashing for blockchain mining and sequence alignment for genomics. The trade-off is total inflexibility: if that core algorithm changes, the ASIC cannot be reprogrammed and becomes obsolete, locking the hardware design to the specific problem version it was built to solve.

The trajectory yields an engineering rule: the era of “free” performance gains from general-purpose scaling is over. For decades, software engineers could rely on Moore’s Law to accelerate existing code without architectural changes. The breakdown of Dennard scaling forced a decisive change: engineers can no longer wait for faster CPUs to solve computational bottlenecks but must instead design the hardware to fit the algorithm. This necessity of hardware-software co-design is why modern AI engineering requires deep understanding of the underlying silicon. Performance is now determined by how well the algorithm’s memory access patterns and parallelism map to the specialized physical structures of domain-specific architectures.

11.2.5 Machine learning hardware specialization

Many important neural-network workloads contain regular tensor operations, substantial data parallelism, and opportunities for reduced precision. Dense matrix multiplications and convolutions are prominent examples because their loop structure exposes repeated operations and predictable reuse. These characteristics have driven specialized hardware architectures that would be ineffective for arbitrary code but can provide substantial gains on matching ML kernels. Other ML operations remain irregular, bandwidth bound, sequential, or too small to fill a wide engine, so the same device need not accelerate every part of a model. The hardware built to exploit the regular portion constitutes a class of devices known as ML accelerators, and the economic trigger for specialization appears when that supported workload is important at fleet scale rather than only in a benchmark.



Example 11.2: The TPU capacity cliff

Scenario: Google projected in 2013 that user adoption of voice-search speech recognition (3 minutes/day/user) would require doubling global data center capacity to run CPU-based neural network inference (Jouppi et al. 2017).

Diagnosis: General-purpose CPU server capacity was economically unable to absorb the exponential growth in inference FLOPs, threatening a capital infrastructure bottleneck.

Systems lesson: Custom hardware acceleration becomes mandatory when workload volume crosses fleet-level economic thresholds. Aggregate arithmetic intensity and server total cost of ownership drive custom ASIC deployment decisions.

Machine learning computational requirements reveal limitations in traditional processors. In an illustrative scenario where a reference CPU sustains 5 percent–10 percent of its peak on a neural network workload, the upper endpoint delivers approximately 0 GFLOP/s (billions of floating-point operations per second). This inefficiency results from architectural mismatches: CPUs optimize for single-thread performance and irregular memory access, while neural networks require massive parallelism and predictable data streams. The memory bandwidth constraint compounds the problem: a single neural network layer may require accessing gigabytes of parameters, overwhelming CPU cache hierarchies designed for smaller working sets.

The energy economics of data movement influence accelerator design. Per event, accessing a 32-bit value from DRAM can consume on the order of $10^2\times$ more energy than one FP32 multiply (exact values vary by technology node and design), making data movement a primary optimization target (Horowitz 2014; Sze et al. 2017). This disparity helps explain the progression from repurposed graphics processors to purpose-built neural network accelerators. TPUs and other custom accelerators can sustain high utilization on dense kernels by implementing systolic arrays and other architectures that maximize data reuse while minimizing movement.

Training and inference present different computational profiles that influence accelerator design. Training computes gradients and updates weights: FP32 and FP16 are standardized binary floating-point formats (IEEE Standards Association 2019b), while mixed-precision training uses lower-precision tensor operations with higher-precision accumulation when accuracy permits (Micikevicius et al. 2017). Backpropagation also retains or recomputes activations (see section 8.3.3.1), increasing memory demand. Inference requires only the forward path and may use INT8 or INT4 when the model and kernels support it. Interactive inference may prioritize tail latency, while offline or batched inference may prioritize throughput, cost, or energy; accelerator choice follows the actual service objective. These profiles often favor high-capacity, throughput-oriented training systems

and energy-efficient inference paths, but batch size and service policy can reverse the apparent preference.¹²



Definition 11.3: ML accelerator

Machine Learning Accelerators are processors or specialized units designed primarily for common neural-network tensor operations and dataflows. They pursue high R_{peak} and data reuse on matching workloads by devoting more resources to arithmetic and local data movement than a general-purpose execution path can.

1. **Significance:** An ML accelerator's defining feature is not raw arithmetic alone but a balanced path that supplies operands and retains reusable values nearby. The A100 provides 2.04 TB/s of device-memory bandwidth, about 10.2× the illustrative 200 GB/s host-DRAM path used here (NVIDIA Corporation 2020a; Choquette et al. 2021). This is a ceiling comparison, not a guarantee: sufficient arithmetic intensity, concurrency, and reuse are still required to approach its 312 TFLOP/s FP16/BF16 peak.
2. **Distinction:** The gains are conditional on operation support, parallel work, regular data access, and a software stack that selects the intended path. An ML accelerator can be orders of magnitude faster than a reference CPU on large dense matrix multiplication, yet lose much of that advantage on small, irregular, synchronization-heavy, or control-heavy workloads. Transfers and unsupported operators can dominate even when the accelerated kernel itself is fast.
3. **Common pitfall:** A frequent misconception is that ML accelerators always accelerate ML. Approaching peak throughput requires a compatible precision, operation shape, layout, implementation, and enough parallel work to keep the relevant units busy. Utilization also depends on whether memory and interconnect paths can sustain that work. A batch-1 autoregressive request may therefore use only a small fraction of a large training accelerator's arithmetic capacity even though the same device performs well on batched prefill or training.

Deployment context shapes architectural choices by identifying the binding constraint. In data centers, the constraint is often time-to-result for training massive models. An NVIDIA H100 trades hundreds of watts of power for high throughput (Choquette 2023); whether that trade lowers total cost depends on utilization, rental price, workload scaling, and the electricity rate. Google's TPUv4 makes a similar architectural trade, prioritizing throughput through systolic arrays and high-bandwidth memory (Jouppi et al. 2023).



Checkpoint 11.2: The accelerator gate

Use the energy hierarchy to decide when specialization pays.

Energy inversion

- ☐ **Data movement cost:** can you explain why moving data from DRAM costs 100× more energy than computing on it?
- ☐ **Architectural response:** explain how systolic arrays (TPU) and Tensor Cores (GPU) reduce repeated movement of the same operands.

Selection logic

- ☐ **Training vs. inference:** why do training chips need massive HBM bandwidth, while inference chips prioritize low latency and INT8 ops?

At the edge, the priority often shifts toward energy per inference and hard latency or power limits, although throughput remains important for continuous camera, audio, or sensor streams. A smartphone camera or always-on audio path operating inside a few-watt budget cannot simply adopt a data center accelerator's high-power memory system. Edge architectures instead reduce

12 | Latency vs. throughput in accelerator design:

Training commonly uses throughput-oriented execution and larger batches to amortize work, while latency-sensitive inference must control single-request and tail latency. This is a service objective, not a universal training/inference divide: batched inference can also be throughput oriented, and small or interactive training loops may care about step latency. Pipeline depth, batching delay, kernel setup, and queueing can make a training-oriented path less suitable for an interactive request even when its peak FLOP/s is higher, so the design must match the workload's actual arrival pattern, batch opportunity, and objective.

DRAM 640 pJ

FP32 mul 3.7 pJ

SRAM 0.5 pJ

log scale

A DRAM access costs over 100× one FP32 multiply per event.

movement through local scratchpads, tightly integrated accelerators, dynamic voltage and frequency scaling, and event-driven processing when the workload permits. The same memory-wall principle applies in both settings: data center chips invest in high-bandwidth memory (HBM) capacity and bandwidth, while edge chips depend heavily on proximity and reuse.

No single architecture dominates every ML workload. Edge devices favor energy efficiency and bounded latency, while cloud-scale training values throughput, capacity, and interconnect performance. Cloud inference can emphasize cost per request or tail latency, and an edge device may still need sustained throughput for a real-time stream. Specialized architectures therefore reflect their deployment context, yet all remain subject to the energy and latency costs of moving data.

Table 11.1 summarizes these milestones in hardware specialization. Floating-point coprocessors accelerated arithmetic previously implemented in software, early GPUs increased graphics throughput, and media engines showed how stable pipelines could justify fixed-function blocks. AI accelerators combine dense tensor units with memory hierarchies and software support, making integration between data movement and parallel execution a central challenge.

Modern AI acceleration therefore extends beyond the chip. Useful accelerators need framework, compiler, library, driver, and runtime support for graph transformations, kernel selection, fusion, memory scheduling, and deployment across environments from data centers to edge devices. Unsupported operations or surrounding transfers can erase a kernel-level gain.

Table 11.1: Hardware Specialization Trends: Successive computing eras progressively integrate specialized hardware to accelerate prevalent workloads, moving from general-purpose CPUs to domain-specific architectures and ultimately to customizable AI accelerators. Tailoring hardware to computational patterns improves performance and energy efficiency, driving innovation in machine learning systems.

Era	Computational pattern	Architecture examples	Characteristics
1980s	Floating-Point & Signal Processing	FPU, DSP	• Single-purpose engines • Focused instruction sets • Coprocessor interfaces
1990s	3D Graphics & Multimedia	GPU, SIMD Units	• Many identical compute units • Regular data patterns • Wide memory interfaces
2000s	Real-time Media Coding	Media Codecs, Network Processors	• Fixed-function pipelines • High throughput processing • Power-performance optimization
2010s	Deep Learning Tensor Operations	TPU, GPU Tensor Cores	• Matrix multiplication units • Massive parallelism • Memory bandwidth optimization
2020s	Application-Specific Acceleration	ML Engines, Smart NICs, Domain Accelerators	• Workload-specific datapaths • Customized memory hierarchies • Application-optimized designs

11.2.6 The integration bottleneck

For many accelerator workloads, adding arithmetic units is easier than feeding them efficiently. Early coprocessors accelerated floating-point arithmetic implemented in software, and GPUs increased throughput for graphics and later general compute. Modern AI systems must also solve an **integration bottleneck**: moving operands and results through a memory hierarchy quickly and efficiently enough to sustain many parallel compute units.

Three properties of many neural-network kernels make this integration tractable. Large matrix multiplications and convolutions expose data parallelism that dense processing-element arrays can exploit. Static or compiled tensor graphs often have analyzable dataflow, allowing compilers to tile transfers into local scratchpads¹³ instead of relying only on hardware-managed caches. This gives software explicit control over selected, regular working sets without eliminating caches or irregular access. Selected operations may also tolerate reduced precision, allowing supported 8-bit or 4-bit paths to increase compute density and reduce bytes per value (Dally et al. 2021; Dally 2023). Dynamic control flow, embeddings, small tensors, variable shapes, and accuracy-sensitive operators fit these assumptions less well and can remain bottlenecks.

Once arithmetic capacity is abundant, the engineering challenge becomes keeping data close enough to use it. A DRAM access can consume more than 100× the energy of a low-precision

¹³ | **Scratchpad memory:** When a compiler can determine a tensor kernel’s access pattern, it can explicitly stage tiles in fast, software-controlled local memory. This avoids some tag, replacement, and coherence machinery, but transfers, synchronization, banking, and capacity must be scheduled correctly. Scratchpads complement rather than universally replace caches because irregular or shared data may still benefit from hardware management. Google’s TPU v1, for example, used a 24 MB software-managed Unified Buffer for intermediate activations, while weights and instructions used other paths (Jouppi et al. 2017).

arithmetic operation on the technology assumptions summarized by Horowitz (2014). This hierarchy helps explain why accelerator designs invest in HBM, on-chip SRAM, and software-managed scratchpads rather than only adding compute units.¹⁴

Hardware acceleration addresses the memory bottleneck through specialized spatial layout and multi-tiered storage. Examine the architectural blueprint in figure 11.5, noting how high-bandwidth memory feeds local scratchpads surrounding the processing element grid.




Figure 11.5: Anatomy of a Modern AI Accelerator: This generic schematic combines common components rather than depicting one product. A host interface connects the accelerator to a CPU, a high-bandwidth device-memory interface feeds shared on-chip storage, and processing elements contain matrix, vector, and special-function units with local memory. Real accelerators differ in which units they provide and whether storage is hardware-managed cache or software-managed scratchpad.

The evolution from the Intel 8087 to the Google TPU reveals a consistent pattern: hardware evolves to fit the algorithm's dominant bottleneck. Where the 8087 addressed floating-point operations that dominated many scientific workloads, modern AI accelerators address dense matrix and convolution operations that dominate much of neural-network training and inference (Palmer 1980; Goodfellow et al. 2016; Sze et al. 2017; Jouppi et al. 2017). This concentration of demand explains why specialized AI silicon can deliver large performance-per-watt improvements over general-purpose processors on matching workloads.

These three opportunities—parallel work, analyzable dataflow, and supported lower precision—shape many accelerator decisions. Their value depends on the accelerator's physical organization and on how well software maps the workload onto it.

In the generic design in figure 11.5, an array of **processing elements** contains units for different operation classes: matrix units execute matrix multiplication, vector units perform element-wise work, and special function units approximate functions such as exponentials. The number, organization, and names of these units vary by architecture, but all seek to exploit data-level parallelism across tensor tiles. Some designs expose an explicit systolic grid; others distribute matrix units among wider parallel cores. The schematic therefore represents functional roles, not a universal physical topology.

The memory hierarchy is equally important. High-bandwidth memory supplies off-chip capacity and aggregate throughput, while shared on-chip storage and per-element caches or scratchpads reduce traffic and energy. The exact hierarchy differs across products: some designs rely more on hardware-managed caches, others expose software-managed SRAM, and many combine both. Moving a tensor tile downward through this hierarchy costs time and energy, so compilers and kernels coordinate capacity, reuse, prefetching, and double buffering. More HBM bandwidth raises the off-chip ceiling but does not remove the need for local reuse. The common purpose is to keep

¹⁴ | **HBM:** The A100 and H100 generations provide 2.0–3.4 TB/s of device-memory bandwidth through stacked memory and wide interfaces, compared with 760 GB/s for the GDDR6X reference used here (NVIDIA Corporation 2020a; Choquette 2023). Raising the bandwidth ceiling can move some kernels toward the compute-bound side of the roofline, but only if their access pattern and arithmetic intensity can use it. HBM also increases packaging and system cost, so its value depends on the workload's bandwidth demand.

reusable weights, activations, and partial results near the units that consume them. The machine foundations appendix collects reference specifications for modern accelerators, including H100 and TPU v5, and summarizes the latency hierarchy.

The host interface connects the accelerator to the broader system. A CPU commonly coordinates program execution, I/O, and unsupported operations, while the accelerator executes the kernels assigned to it. This partition is not absolute: accelerators can schedule work locally, and CPUs can execute tensor operations. What matters is where each operator runs and which transfers, launches, and synchronization points surround it. Direct-memory-access engines and device-side work queues can overlap some of that overhead with computation, but unsupported operators may still force costly handoffs. The path from host through device memory and on-chip storage to the processing elements therefore defines the useful boundary of acceleration. A fast kernel does not yield a fast application when communication across that boundary dominates end-to-end time.

With the accelerator’s physical architecture established, the next step is to explain why these specific components dominate. Tensor Cores, vector units, and hierarchical memory do not exist by accident; they exist because neural network computations repeatedly invoke a small set of operations. Understanding these patterns is essential because they explain which algorithmic changes translate to real speedups (those that align with hardware primitives) and which remain purely theoretical.

11.3 AI Compute Primitives

Across fully connected, convolutional, and attention layers, learned linear transformations reduce to repeated multiply-accumulate (MAC) operations. A large layer can issue millions or billions of MACs, and its regular tensor structure exposes parallel work and reuse, making it a natural specialization target. MACs do not dominate every model or phase—normalization, routing, indexing, communication, and element-wise operations can become limiting—but they account for much of the peak arithmetic capacity in modern ML accelerators. Whether they dominate elapsed time still depends on tensor shape, data movement, and the surrounding operators.

The hardware units that exploit these patterns are **AI compute primitives**: specialized functional blocks, each optimized for a particular class of operation. Three primitives are especially common in accelerators, each targeting a distinct computational pattern found in neural networks.

Listing 11.1 demonstrates how a dense layer decomposes at the framework level, encapsulating thousands of multiply-accumulate operations in a single high-level call.

Listing 11.1: Dense Layer Abstraction: High-level framework APIs encapsulate 131,072 MACs (256 inputs times 512 outputs) in a single function call, hiding the computational complexity from developers while enabling automatic hardware optimization.

```
# Framework abstracts compute-intensive operations
dense = Dense(512)(input_tensor) # 256x512 MACs per sample
```

This single line of code conceals the computational complexity that accelerators must handle. Listing 11.2 reveals how the framework expands this high-level call into mathematical operations.

Listing 11.2: Matrix Operation Expansion: Each dense layer decomposes into matrix multiplication and element-wise operations, exposing the dominant compute pattern that many neural-network kernels are built around.

```
# Linear transformation work scales with input_dim x output_dim x
# batch.
output = (
    matmul(input, weights) + bias
) # Matrix multiply dominates cost
output = activation(
    output
) # Element-wise: proportional to output_dim x batch
```

The matrix multiplication dominates computation time, but this abstraction still hides the underlying loop structure. At the processor level, listing 11.3 reveals how nested loops multiply inputs and

weights, sum the results, and apply a nonlinear function, exposing the $\mathcal{O}(B \times d_{\text{in}} \times d_{\text{out}})$ complexity that accelerators must handle efficiently.

Listing 11.3: Processor-Level Execution: Nested loops reveal the $\mathcal{O}(B \times d_{\text{in}} \times d_{\text{out}})$ multiply-accumulate operations that accelerators must execute, with 4.2M MACs for $B=32$, $d_{\text{in}}=256$, $d_{\text{out}}=512$ configurations.

```
# Total operations: batch_size * output_size * input_size MACs
for n in range(batch_size): # Batch dimension: parallelizable
    for m in range(output_size): # Output neurons: parallelizable
        sum = bias[m] # Initialize accumulator
        for k in range(input_size): # Reduction dimension: sequential
            sum += input[n, k] * weights[k, m] # MAC operation
        output[n, m] = activation(sum) # Nonlinear transformation
# Example work scales as batch_size * output_size *
# input_size multiply-accumulate operations
```

This loop structure reveals three computational patterns that recur across many neural-network architectures: element-wise operations along vectors, matrix-level reductions, and nonlinear transformations. Their frequency justifies specialized execution paths, although the gain depends on tensor shape, data movement, and operator support. These patterns become hardware blocks: vector units handle independent elements and reductions, matrix engines organize multiply-accumulate work across tiles, and special-function paths approximate nonlinear math. The boundaries are architectural rather than mathematical; one fused kernel may use several paths without exposing the transitions to the framework.

11.3.1 Vector operations

Vector operations provide one level of acceleration by processing multiple data elements per instruction. Recall the nested-loop structure exposed in listing 11.3: a batch of 32 samples through a 256-to-512 dense layer requires 4.2M MACs. A scalar formulation loads one input-weight pair, multiplies it, and updates one accumulator at a time. Vector execution performs that work across several lanes, amortizing instruction decoding, loop control, and address generation. Its advantage depends on vector width, alignment, tail handling, and enough memory bandwidth to supply the lanes.

Listing 11.4: Vectorized Dot-Product Loop: This illustrative loop processes several element pairs per iteration and reduces them into one scalar dot product.

```
fmv.w.x fa0, zero # Scalar dot-product accumulator
fmv.w.x ft1, zero # Zero seed for vector reduction
loop_feature:
    vsetvli t0, feature_cnt, e32, m1, ta, ma
    vle32.v v1, (in_ptr)
    vle32.v v2, (wt_ptr)
    vmul.vv v3, v1, v2
    fmv.v.f v0, ft1
    vfredusum.vs v4, v3, v0
    fmv.f.s ft0, v4
    fadd.s fa0, fa0, ft0
    slli t1, t0, 2 # Four bytes per FP32 element
    add in_ptr, in_ptr, t1
    add wt_ptr, wt_ptr, t1
    sub feature_cnt, feature_cnt, t0
    bnez feature_cnt, loop_feature
```

Vector processing units solve this by operating on multiple data elements simultaneously. RISC-V,¹⁵ an open instruction-set architecture (ISA) with a standardized vector extension (Waterman et al. 2013), provides a useful setting for illustrating this idea. Listing 11.4 uses vector-style assembly code in which a single instruction processes several data elements. The loop has five hardware-visible stages:

¹⁵ | **RISC-V (reduced instruction set computer V):** The open ISA allows hardware teams to add custom ML instructions, including vector dot products, activation functions, and sparse tensor operations, without changing the base ISA. The trade-off is software ecosystem maturity: custom extensions require corresponding compiler, library, and runtime support, which can limit portability and increase integration work.

1. **Vector length configuration:** Configures the vector units to process 32-bit elements, automatically determining how many operations happen in parallel based on hardware width (VLEN).
2. **Vector initialization:** Clears the scalar accumulator that will hold the completed dot product.
3. **Vector loads:** Loads contiguous 32-bit values into v1 and v2, using one vector-load instruction for each source.
4. **Vector multiply and reduction:** Multiplies corresponding elements and reduces the products to a partial scalar sum.
5. **Pointer arithmetic:** Advances the pointers and decrements the remaining element count by the run-time vector length.

In this assembly sequence, the vector multiply instruction processes several element pairs at once, while the reduction completes each partial dot product and the vector loads amortize instruction overhead across multiple data elements. The run-time vector length determines how many values each iteration transfers and processes, so the same loop works across implementations with different vector widths. The illustrative workload is divided into roughly 524,288 vector chunks, each executed by several vector and scalar instructions.

Key vector operations map directly to common deep learning patterns. Table 11.2 enumerates how operations such as reduction, gather, scatter, and masked operations appear frequently in pooling, embedding lookups, and attention mechanisms, clarifying the direct mapping between low-level vector hardware and high-level machine learning workloads.

Table 11.2: Vector Operations: Core vector operations map directly to deep learning primitives: reductions implement pooling layers, gathers enable embedding lookups, scatters update embedding gradients, and masked operations handle attention masks. These mappings show how vector hardware supports common ML primitives.

Vector operation	Description	Neural network application
Reduction	Combines elements across a vector (for example, sum, max)	Pooling layers, attention score computation
Gather	Loads multiple nonconsecutive memory elements	Embedding lookups, sparse operations
Scatter	Writes to multiple nonconsecutive memory locations	Gradient updates for embeddings
Masked operations	Selectively operates on vector elements	Attention masks, padding handling
Vector-scalar broadcast	Applies scalar to all vector elements	Bias addition, scaling operations

The benefit extends beyond instruction count. Contiguous vector loads can use memory interfaces efficiently, and control overhead is amortized across several data elements. Gather, scatter, masking, and short tails may use fewer lanes or require extra memory transactions, so vector width alone does not guarantee proportional speedup. The architectural pattern is not new. The Cray-1¹⁶ used vector registers and pipelined functional units for scientific computing in the 1970s (Jordan 1982); modern ML processors apply related principles at much larger commercial scale.

Vector operations excel at element-wise transformations like activation functions, where each output depends only on *its corresponding input*. Neural networks, however, also require structured computations where each output depends on *all inputs*—the weighted sums that define layer transformations. These many-to-many operations naturally express themselves as matrix multiplications, our second compute primitive.

11.3.2 Matrix operations

Matrix multiplication accounts for much of the arithmetic in many dense neural networks, transforming high-dimensional data through structured patterns of weights, activations, and gradients (Goodfellow et al. 2016). While vector operations process elements independently or reduce them, matrix operations organize work across multiple dimensions. Hardware and libraries tile those

¹⁶ | **Cray-1 vector legacy:** The Cray-1 (1975) used 64-element vector registers with pipelined functional units, allowing a stream of elements to advance through arithmetic operations without issuing one scalar instruction per element. Modern accelerators extend the same principles of wide registers, pipelined execution, and data reuse to vector and matrix tiles.

dimensions so blocks of weights and activations can be reused while partial sums accumulate locally. This regular structure drives important hardware strategies, but small or irregular matrices may leave tile capacity unused.

11.3.2.1 Matrix operations in neural networks

Neural network computations decompose into hierarchical matrix operations. Listing 11.5 captures this hierarchy through a linear layer that transforms input features into output neurons over a batch.

Listing 11.5: Linear-Layer Operation Hierarchy: The example expands a linear layer into a matrix multiply and bias broadcast, then applies element-wise ReLU, separating the matrix and vector work that hardware must execute.

```
layer = nn.Linear(256, 512) # Layer transforms 256 inputs to 512 outputs
output = layer(input_batch) # Process a batch of 32 samples

# Framework Internal: Core operations (column-batch convention)
Z = matmul(weights, input) # Matrix: transforms [256x32]
# input to [512x32] output
Z = Z + bias # Vector: adds bias to each
# output independently
output = relu(Z) # Vector: applies activation to
# each element independently
```

This computation demonstrates the scale of matrix operations in neural networks. Each output neuron (512 total) must process all input features (256 total) for every sample in the batch (32 samples). The weight matrix alone contains $256 \times 512 = 131,072$ parameters that define these transformations, illustrating why efficient matrix multiplication dominates performance considerations.

Neural networks employ matrix operations across architectural patterns beyond simple linear layers. Convolution operations transform into matrix multiplications through the im2col technique,¹⁷ enabling efficient execution on matrix-optimized hardware. Listing 11.6 illustrates these applications.

Listing 11.6: Matrix Patterns Across Architectures: Linear layers, attention mechanisms, and convolutions all reduce key work to matrix multiplications, making matrix hardware the shared primitive across modern neural architectures.

```
hidden = matmul(weights, inputs)
# weights: [out_dim x in_dim], inputs: [in_dim x batch]
# Result combines all inputs for each output

# Attention Mechanisms - Multiple matrix operations
Q = matmul(Wq, inputs)
# Project inputs to query space [query_dim x batch]
K = matmul(Wk, inputs)
# Project inputs to key space [key_dim x batch]
attention = matmul(Q.T, K)
# Compare all query tokens with all key tokens [batch x batch]

# Convolutions - Matrix multiply after reshaping
patches = im2col(input)
# Convert [H x W x C] image to matrix of patches
output = matmul(kernel, patches)
# Apply kernels to all patches simultaneously
```

These examples differ in tensor shape and reuse pattern, which determines how efficiently each multiplication occupies a matrix unit. Linear layers reuse one weight matrix across a batch; attention forms projections and then a token-to-token score matrix; convolution reuses each kernel across many image patches. Expressing all three as matrix multiplication exposes a common interface to hardware, but it does not make their execution costs identical. Batch size, sequence length, channel count, and tiling determine whether operands are reused locally or repeatedly fetched from memory.

¹⁷ | **Im2col (image-to-column):** Transforms convolution into a matrix multiplication by arranging receptive-field values as columns. An explicit materialization can duplicate overlapping values and expand storage by up to roughly the kernel area before edge, stride, and channel effects are considered. Production libraries often use implicit-GEMM or direct-convolution kernels to obtain matrix-friendly execution without materializing the full expanded matrix.

11.3.2.2 Matrix operations hardware acceleration

This pervasive pattern of matrix multiplication has direct implications for hardware design: accelerators need specialized units that can handle these computations at scale. Listing 11.7 demonstrates a representative dedicated matrix unit that processes an entire 16×16 block at once, illustrating why matrix instructions and Tensor Cores can deliver much higher throughput than scalar or vector-only execution paths (NVIDIA 2017; Intel Corporation 2021a).

Listing 11.7: Illustrative Matrix-Unit Operation: Representative load, matrix-multiply-accumulate, and store instructions operate on matrix blocks rather than individual scalar elements.

```
mload mr1, (weight_ptr)    # Load e.g., 16x16 block of
                             # weight matrix
mload mr2, (input_ptr)     # Load corresponding input block
matmul.mm mr3, mr1, mr2    # Multiply and accumulate entire
                             # blocks at once
mstore (output_ptr), mr3   # Store computed output block
```

The illustrative unit operates on matrix tiles rather than individual elements. A full $16 \times 16 \times 16$ tile product contains $16^3 = 4,096$ multiply-accumulate operations, but the cycles required depend on the actual instruction and architecture. Sustained throughput also depends on the number of matrix units, operand precision, issue rate, pipeline behavior, and whether memory can keep the units fed. Matrix units complement vector execution by accelerating structured many-to-many transformations.

Like vector processing, matrix acceleration has deep historical roots—DSPs and GPUs optimized for matrix computations in the 1980s-1990s for image processing, scientific computing, and 3D rendering (Owens et al. 2008; Hwu 2011). Neural networks have made matrix multiplication commercially dominant, driving the development of dedicated Tensor Cores and TPUs that process these operations at unprecedented scale.

Matrix and vector operations together handle the linear algebra of neural networks. Between every linear transformation, however, sits a nonlinear activation function—and these transcendental computations (exponentials, square roots, trigonometric functions) cannot be efficiently expressed through multiply-accumulate alone. Table 11.3 summarizes the three execution roles, clarifying which neural network operations map to each.

Table 11.3: Operation Characteristics: Matrix operations excel at many-to-many transformations, while vector units handle element-wise operations and reductions such as activation and normalization. Special-function paths handle transcendental arithmetic through architecture-specific approximation or dedicated logic. The operation structure determines which hardware primitive is appropriate.

Operation type	Best For	Examples	Key characteristic
Matrix Operations	Many-to-many transforms	Layer transformations, attention, convolutions	Each output depends on multiple inputs
Vector Operations	Vector and reduction work	Activation functions, layer normalization, element-wise gradients	Uses element-wise operations and vector reductions
Special-Function Paths	Transcendental arithmetic	Exponential, logarithm, reciprocal square root	Uses architecture-specific approximation or dedicated logic

11.3.3 Special function units

Special Function Units (SFUs) or specialized instruction paths accelerate nonlinear functions and related arithmetic, completing the chapter’s trio of processing primitives. The need is not new: floating-point coprocessors addressed scalar arithmetic bottlenecks (Palmer 1980), and digital signal processors added specialized arithmetic for signal-processing workloads (Smith 1997). An implementation may use a standalone unit, vector approximation instructions, or a library sequence, trading accuracy, latency, and throughput. In neural networks, activation, normalization, and softmax operations can become important between matrix kernels, particularly when repeated memory passes or transcendental functions limit throughput.

11.3.3.1 Nonlinear functions

To see why dedicated hardware matters, consider a typical layer sequence (Goodfellow et al. 2016). Listing 11.8 combines linear transformations with nonlinear activations—operations that appear simple in Python but reveal substantial computational complexity at the hardware level.

Listing 11.8: Linear–Nonlinear Layer Sequence: A linear layer, ReLU, and BatchNorm combine matrix, element-wise, and normalization work in one common layer sequence.

```
layer = nn.Sequential(
    nn.Linear(256, 512), nn.ReLU(), nn.BatchNorm1d(512)
)
output = layer(input_tensor)
```

This sequence introduces multiple nonlinear transformations that extend beyond simple matrix operations. Listing 11.9 breaks down these operations into their mathematical components, exposing the computational complexity that hardware must address.

Listing 11.9: Nonlinear Operation Expansion: Expanding the layer sequence exposes ReLU comparison, BatchNorm reductions, variance computation, and square-root normalization around the matrix multiply.

```
Z = matmul(weights, input) + bias # Linear transformation
H = max(0, Z) # ReLU activation
mean = reduce_mean(H, axis=0) # BatchNorm statistics
var = reduce_mean((H - mean) ** 2) # Variance computation
output = gamma * (H - mean) / sqrt(var + eps) + beta # Normalization
```

11.3.3.2 Hardware implementation of nonlinear functions

The computational complexity of these operations becomes apparent when examining their implementation on traditional processors. These seemingly simple mathematical operations translate into complex sequences of instructions. Consider batch normalization (Ioffe and Szegedy 2015): computing the normalization requires reductions, variance calculation, and a square root, while operations like softmax introduce exponentials whose cost depends on the processor implementation. A rectified linear unit (ReLU) is mathematically simple, but a naive scalar implementation still performs a comparison and selection for every element; optimized ML kernels usually make that step branchless. Listing 11.10 therefore uses ReLU and batch normalization to show two different sources of overhead: element-wise passes through memory and multi-pass normalization work.

The label *nonlinear* hides three different execution contracts. Pointwise functions such as ReLU leave elements independent and therefore map naturally to vector lanes or parallel threads. Reductions such as the mean and variance in batch normalization introduce cross-element dependencies: partial results must be combined before later stages can proceed. Transcendental functions such as exponential, logarithm, and reciprocal square root add a numerical contract because their implementations trade latency, precision, and valid input range. The dependency pattern, not the mathematical label alone, determines whether vector arithmetic, a reduction tree, or a specialized approximation path is appropriate. A faster special-function path accelerates only the operation it implements; it does not remove surrounding tensor traffic or synchronization.

Listing 11.10 is intentionally dependency-explicit rather than a model of an optimized kernel. The loops expose three state lifetimes: per-element activations, per-feature accumulators, and the final normalized outputs. The normalization parameters cannot be applied until the required statistics are available, so implementations must either retain intermediate activations or recompute them. Training derives those statistics from the current batch, whereas inference can use stored statistics and reduce normalization to an element-wise affine transform. Compilers may fuse legal stages, but the data dependencies determine which transformations preserve the computation. Reading the listing as a dependency graph therefore reveals the durable hardware questions: how often each tensor crosses a memory interface, which intermediates remain local, and where reductions impose synchronization.

Listing 11.10: ReLU and BatchNorm Operations: Neural networks process input data through element-wise selections and multiple normalization passes, highlighting efficiency challenges in naive implementations.

```

for batch in range(32):
    for feature in range(512):
        # ReLU: Naive scalar compare/select; optimized kernels
        # usually implement this branchlessly.
        z = matmul_output[batch, feature]
        h = max(0.0, z) # Conditional operation

        # BatchNorm: Multiple passes over data
        mean_sum[feature] += h # First pass for mean
        var_sum[feature] += h * h # Additional pass for variance

        temp[batch, feature] = h # Extra memory storage needed

# Normalization requires complex arithmetic
for feature in range(512):
    mean = mean_sum[feature] / batch_size
    var = (var_sum[feature] / batch_size) - mean * mean

    # Square root computation: Multiple iterations
    scale = gamma[feature] / sqrt(var + eps) # Iterative
    # approximation
    shift = beta[feature] - mean * scale

    # Additional pass over data for final computation
    for batch in range(32):
        output[batch, feature] = temp[batch, feature] * scale + shift

```

Each non-linear operation introduces distinct hardware challenges across deep network layers. Batch normalization requires multiple passes through data: one for mean computation, another for variance, and a final pass for output transformation. Each pass loads and stores data through the memory hierarchy. Operations that appear simple in mathematical notation often expand into many instructions, especially for square roots and exponentials on processors without specialized hardware paths. ReLU generally maps to a compare-and-select or maximum operation, so its standalone cost is dominated less by arithmetic than by the additional read and write if it is not fused with neighboring work. The implementation needs temporary storage for intermediate values, increasing memory usage and bandwidth consumption. While vector units excel at regular computations, functions like exponentials and square roots often require specialized implementations that may not fully use vector processing capabilities.

11.3.3.3 SFU hardware implementation

SFUs address these inefficiencies through dedicated hardware implementation. Modern ML accelerators include specialized circuits that transform these complex operations into low-latency, fixed-function computations. Listing 11.11 illustrates the mapping: after one vector load, architecture-dependent ReLU, sigmoid, tanh, and reciprocal-square-root paths consume the same input vector.

Listing 11.11: Illustrative SFU Operations: Pseudocode shows vector nonlinear operations that may map to specialized execution paths; instruction names and latency are implementation dependent.

```

vld.v v1, (input_ptr)    # Load vector of values (pseudocode)
vrelu.v v2, v1           # Vector ReLU path
vsigm.v v3, v1           # Vector sigmoid path
vtanh.v v4, v1           # Vector tanh path
vrsqrt.v v5, v1          # Vector reciprocal-square-root path

```

Specialized function paths implement nonlinear and reduction primitives through architecture-specific circuitry. A ReLU can use compare-and-select logic; square root, exponential, and logarithmic

functions may use iterative approximations, lookup tables, or interpolation. Table 11.4 summarizes representative mechanisms and their qualitative latency behavior.

Table 11.4: Special Function Units: Dedicated paths reduce instruction overhead for nonlinear and reduction primitives. Implementations and latency are architecture-dependent; the table describes mechanisms rather than universal cycle counts.

Function unit	Operation	Implementation strategy	Illustrative latency
Activation unit	ReLU	Compare-and-select or maximum	Low; architecture-dependent
Statistics unit	Mean, variance	Parallel reduction trees	Grows with reduction depth
Exponential unit	exp, log, sigmoid, tanh	Approximation, table lookup, and interpolation	Architecture-dependent
Root/power unit	sqrt, rsqrt	Iterative or approximation hardware	Architecture-dependent

Vector operations, matrix operations, and special-function paths cover three important computational patterns, but primitives alone do not determine throughput. The primitives tell us *what* an accelerator can execute efficiently; the execution model and memory system determine *how* the work maps onto parallel hardware. The same matrix multiplication can therefore achieve very different fractions of peak performance as tensor shape, tile alignment, thread organization, occupancy, memory access, and synchronization change. Peak ratings describe a supported operation under favorable conditions; execution models explain how much of that capacity a real workload can sustain.

11.4 Compute Units and Execution Models

Applying ReLU to a 512-element vector shows why execution models matter: the operation is simple, but throughput depends on whether the hardware treats those 512 comparisons as scalar instructions, SIMD lanes, GPU threads, or tensor-program fragments. Modern AI processors package the three compute primitives into distinct execution units: single instruction, multiple data (SIMD) units, Tensor Cores, and processing elements that define how computations are structured and exposed to programmers. Understanding this organization reveals both the theoretical capabilities and practical performance characteristics that determine real-world throughput.

11.4.1 Mapping primitives to execution units

The progression from computational primitives to execution units follows a structured hierarchy that reflects the increasing complexity and specialization of AI accelerators:

- Vector operations → SIMD and single instruction, multiple threads (SIMT) units that enable parallel processing of independent data elements
- Matrix operations → Tensor Cores¹⁸ and systolic arrays that provide structured matrix multiplication
- Special functions → Dedicated hardware units integrated within processing elements

Each execution unit combines these computational primitives with specialized memory and control mechanisms, optimizing both performance and energy efficiency. This structured packaging allows hardware vendors to expose standardized programming interfaces while implementing diverse underlying architectures tailored to specific workload requirements. The choice of execution unit significantly influences overall system efficiency by determining data locality, compute density, synchronization overhead, and how much of the theoretical peak the workload can actually use.

11.4.2 Evolution from SIMD to SIMT architectures

Imagine applying ReLU to a 512-element vector. A scalar formulation contains 512 comparison-and-select operations. A SIMD instruction might cover 8 or 16 elements at a time, while a **SIMT** GPU can express one lightweight thread per element and schedule those threads in warps or waves. Actual instruction counts and throughput depend on the processor, compiler, memory path, and surrounding fusion. The progression illustrates two related ideas: Flynn’s SIMD taxonomy

¹⁸ | **Reduced-precision ML:** Halving operand width halves the bytes per stored element and can support denser arithmetic, although realized throughput depends on the available datapath and kernel (Dally et al. 2021; Dally 2023). NVIDIA’s transition from Pascal P100 vector arithmetic to mixed-precision Tensor Cores in Volta V100 increased advertised peak FP16 throughput from about 21.2 to 125 TFLOP/s, roughly 6× across those products (NVIDIA 2017). Precision selection must preserve numerical behavior and use an efficient supported path; the smallest format is not automatically the best one.

formalized data-parallel execution (Flynn 1966), and GPU SIMT architectures expose many logical threads that execute in hardware groups (Lindholm et al. 2008; Nickolls et al. 2008).

SIMD execution applies identical operations to multiple data elements in parallel, minimizing instruction overhead while maximizing data throughput. This execution model is widely used to accelerate workloads with regular, independent data parallelism, such as neural network computations. The Arm Scalable Vector Extension (SVE) provides a representative example of how modern architectures implement scalable SIMD operations efficiently (Stephens et al. 2017). Listing 11.12 demonstrates this approach.

Listing 11.12: Arm SVE Vector Execution: Predicated vector assembly loads, multiplies, and adds floating-point elements using length-agnostic registers.

```
ptrue p0.s           # Create predicate for vector length
ld1w z0.s, p0/z, [x0] # Load vector of inputs
fmul z1.s, z0.s, z0.s  # Multiply elements
fadd z2.s, z1.s, z0.s  # Add elements
st1w z2.s, p0, [x1]    # Store results
```

The `ptrue` instruction activates all available predicate lanes without encoding a fixed vector length. This vector-length-agnostic programming model allows the same SVE binary to run on implementations from 128 to 2048 bits without recompilation (Stephens et al. 2017). Intel’s Advanced Matrix Extensions (AMX) are a different kind of specialization: tile registers and matrix instructions expose two-dimensional matrix operations directly to software rather than merely widening a vector lane (Intel Corporation 2021a). Together, SVE and AMX show two hardware-facing paths for ML kernels: vector-length-portable SIMD and fixed tile-based matrix acceleration.

To address these limitations, SIMT extends SIMD principles by enabling parallel execution across multiple independent threads, each maintaining its own program counter and architectural state (Lindholm et al. 2008; Nickolls et al. 2008). This model maps naturally to matrix computations, where each thread processes different portions of a workload while still benefiting from shared instruction execution. In NVIDIA’s GPU architectures, each Streaming Multiprocessor (SM)¹⁹ coordinates thousands of threads executing in parallel, allowing for efficient scaling of neural network computations. Threads are organized into warps,²⁰ which are the basic execution units that enable SIMT efficiency. Listing 11.13 shows this parallel processing model in action.

Listing 11.13: SIMT Execution: Each thread processes a unique output element in parallel, demonstrating how SIMT enables efficient matrix multiplication on GPUs.

```
__global__ void matrix_multiply(float* C, float* A, float*
                                B, int N) { // CUDA kernel
    // Each thread processes one output element
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;

    float sum = 0.0f;
    for (int k = 0; k < N; k++) {
        // Threads in a warp execute in parallel
        sum += A[row * N + k] * B[k * N + col];
    }
    C[row * N + col] = sum;
}
```

The CUDA²¹ kernel assigns one output element to each thread, exposing independent work that the GPU can schedule across many warps. Similar execution models appear in AMD’s RDNA and Intel’s Xe architectures, reinforcing SIMT as a core mechanism for AI acceleration.

¹⁹ | **SM (streaming multiprocessor):** The physical hardware engine that implements the SIMT model by using warp schedulers to coordinate many parallel threads. Maintaining enough active warps can help hide instruction and memory latency, but occupancy alone does not identify the bottleneck. Low occupancy may result from register use, shared-memory use, block dimensions, or insufficient parallel work, while a memory-bound kernel can still have high occupancy.

²⁰ | **Warp:** On NVIDIA GPUs, a warp is a hardware scheduling group of 32 threads that execute a common instruction when their control paths agree. If threads take different paths, the hardware executes the required paths under different masks, reducing useful lane utilization. The penalty depends on path length and how many lanes take each path, which is why ML kernels often prefer uniform or predicated control flow.

²¹ | **CUDA (compute unified device architecture):** Released by NVIDIA in 2006, CUDA eliminated the need to disguise general-purpose computations as graphics operations, opening GPUs to scientific and ML workloads through a C-like programming model. The ecosystem it created—cuBLAS, cuDNN, TensorRT—constitutes a software moat that can lock the ML training stack to NVIDIA hardware: migrating away requires rewriting or replacing thousands of GPU-optimized kernels, a cost that often exceeds the hardware savings of competing platforms. This software lock-in, not raw silicon performance alone, helps explain why many large ML training stacks remain CUDA-centered.

11.4.3 Tensor Cores

Consider a single transformer attention head computing the $\mathbf{Q} \times \mathbf{K}^T$ product for a 2,048-token sequence with 64-dimensional embeddings. This operation requires multiplying a 2048×64 matrix by a 64×2048 matrix: roughly 268.4M MACs, or about 536.9 MFLOP when the multiply and add are counted separately. On a scalar processor executing one FLOP per cycle at 2 GHz, this single attention head would take about 268.4 ms. GPU SIMT execution can distribute this work across many threads, but Tensor Cores go further by processing entire matrix tiles per instruction; under the illustrative assumption used here, the tiled path completes the same operation in 0.5 milliseconds, a roughly $536.9\times$ improvement over scalar execution. This dramatic speedup arises not from faster clock speeds but from a fundamentally different approach to organizing computation around matrix blocks rather than individual elements.

While SIMD and SIMT units execute vector work efficiently, large matrix computations benefit from units organized around multidimensional tiles. Dedicated matrix engines amortize operand loads across many multiply-accumulate operations by reusing tiles on chip. NVIDIA GPUs expose Tensor Cores, whereas Google TPUs use matrix units built around systolic arrays. Both execute tiled matrix multiplication and accumulation.²² The benefit depends on tile reuse, shape, precision, and memory behavior.

Tensor Cores²³ provide an example of this approach. Listing 11.14 exposes matrix computation capabilities through specialized instructions that use dedicated hardware blocks.

Listing 11.14: Tensor Core Operation: Matrix multiplications are performed in parallel across entire matrix blocks, optimizing computational efficiency for neural network training.

```
Tensor Core Operation (example GPU PTX):
mma.sync.aligned.m16n8k16.row.col.f32.f16.f16.f32
  {d0,d1,d2,d3},      // Destination registers
  {a0,a1,a2,a3},      // Source matrix A
  {b0,b1},            // Source matrix B
  {c0,c1,c2,c3};      // Accumulator
```

A single Tensor Core instruction processes an entire matrix block while maintaining intermediate results in local registers, improving computational efficiency compared to implementations based on scalar or vector operations. This structured approach enables hardware to achieve high throughput while reducing the burden of explicit loop unrolling and data management at the software level.

Design priorities determine how matrix engines appear in different processor families. GPU Tensor Cores preserve programmability while accelerating general-purpose deep learning kernels. TPU-style designs use large-scale matrix units arranged in systolic arrays to maximize sustained training throughput on dense tensor kernels. Mobile NPUs²⁴ shrink the same idea into low-power inference blocks, while server CPUs add matrix instruction extensions (AMX-class tiles) for inference and mixed workloads. Each version changes the same contract: how much flexibility the hardware keeps while reducing movement around dense matrix operations.

Figure 11.6 shows how specialization, lower precision, and sparsity support changed advertised peak capability. From K20X FP32 to H100 FP8 with structured sparsity, the displayed values rise by roughly three orders of magnitude (NVIDIA Corporation 2017, 2020a, 2024; Choquette 2023). The points do not measure one consistent quantity: they mix precision formats and combine FLOP/s with INT8 operations per second. The curve therefore compares each generation's most prominent advertised path, not same-precision application performance. Table E.1 provides the corresponding precision-specific specifications, bandwidth, and power limits.

11.4.4 Processing elements

The highest level of execution unit organization integrates multiple Tensor Cores with local memory into processing elements (PEs). A processing element serves as the primary building block in many AI accelerators, combining different computational units to efficiently execute neural network operations. Each PE typically includes vector units for element-wise operations, Tensor Cores for

²² | **Tensor Core dimension alignment:** NVIDIA Tensor Cores are most efficient when matrix dimensions are aligned to precision- and architecture-specific multiples, such as multiples of 8 or 16 for common FP16/BF16/INT8 paths. Modern cuBLAS and cuDNN can still use Tensor Cores for many nonaligned dimensions, but poorly aligned shapes may trigger less efficient kernels, require padding, or reduce effective throughput (NVIDIA 2024a; NVIDIA Corporation 2021b). This is why model architects often choose embedding and channel dimensions such as 512 rather than 500 and why batch-size-1 inference may fail to reach peak Tensor Core utilization: alignment and arithmetic intensity jointly determine whether the hardware's matrix engines stay full.

²³ | **Tensor Core:** A single Tensor Core instruction executes a complete matrix-multiply-accumulate operation on a small tile of data using a dedicated hardware block (NVIDIA 2017; NVIDIA Corporation 2020a). This approach bypasses the overhead of fetching and scheduling dozens of individual arithmetic instructions on general-purpose CUDA cores. Because these blocks constitute a large fraction of a modern accelerator's advertised tensor throughput, failing to use them can leave most of the chip's theoretical peak unavailable to the workload.

²⁴ | **NPU (neural processing unit):** Mobile NPUs achieve low-power inference by implementing common tensor operations in fixed-function or narrowly programmable hardware rather than as fully general GPU kernels. This architectural commitment can deliver large energy-efficiency gains for supported kernels, but it makes deployment dependent on operator coverage: unsupported functions must fall back to a CPU or GPU path that may be far less efficient for that workload (Sze et al. 2017).




Figure 11.6: Advertised GPU Peak Capability Across Generations: Single-chip values are plotted on a logarithmic axis. The precision label at each point is essential: the curve combines architectural gains with different numerical formats and, for some points, sparsity assumptions rather than measuring same-precision application scaling (NVIDIA Corporation 2017, 2020a, 2024; Choquette 2023).

matrix computation, special function units for nonlinear transformations, and dedicated memory resources to optimize data locality and minimize data movement overhead.

Processing element design varies because each architecture chooses a different balance between compute density, local memory, and interconnect distance. Graphcore’s Intelligence Processing Unit (IPU) distributes computation across 1,472 tiles, each containing independent processing elements optimized for fine-grained parallelism (Graphcore 2020). Cerebras extends the same local-compute principle in the CS-2 system, integrating roughly 850,000 AI-optimized cores across a wafer-scale device for deep learning acceleration (Systems 2021). Tesla’s D1 processor emphasizes substantial local memory inside its processing elements, optimizing throughput and latency for real-time autonomous vehicle workloads (Tesla, Inc. 2021).

Across these designs, the binding trade-off is the one the roster illustrates: compute density versus data locality. Packing more cores raises peak throughput only if each one can be kept fed, so a processing element’s delivered efficiency depends as much on interconnect strategy and memory locality as on raw arithmetic capability.

That same dependence on locality governs which algorithmic optimizations the hardware can actually exploit. A regular grid of processing elements accelerates sparsity only when the surviving nonzero values preserve the predictable access patterns the grid depends on, which is precisely the constraint that N:M structured sparsity is designed to satisfy.

11.4.5 N:M structured sparsity mechanics

While unstructured pruning reduces model size, it rarely translates to hardware speedup because memory access becomes irregular. Hardware accelerators solve this with N:M structured sparsity,²⁵ a pattern-based approach that enforces regularity. The notation “N:M” specifies that exactly N values must be nonzero within every contiguous block of M values, creating a predictable pattern that hardware can exploit.

NVIDIA’s Sparse Tensor Cores implement a concrete instance of this pattern: the 2:4 constraint, which requires that exactly two of every contiguous block of four values be nonzero (equivalently, two must be zero) (NVIDIA Corporation 2020a; NVIDIA 2020a). This constraint allows the hardware to compress the matrix by 50 percent in memory plus metadata. The execution proceeds in three stages: first, the hardware stores only the two nonzero values and compact metadata for every four-

²⁵ | **N:M structured sparsity:** The 2:4 ratio (50 percent density) used by NVIDIA’s Ampere Sparse Tensor Cores is a hardware-friendly compromise: every contiguous four-value group retains two nonzero values, preserving regular indexing while halving the dense value payload (NVIDIA Corporation 2020a). At 2:4, the metadata overhead is compact enough to store alongside the weights without overwhelming the memory-traffic savings, which is the constraint that makes the advertised $2\times$ tensor-math throughput path plausible when kernels and model weights satisfy the pattern.

element block (compression); second, during matrix multiplication, the Sparse Tensor Core reads the metadata to select the corresponding activations and performs math only on the nonzero weights (compute); third, this increases the effective FLOP/byte ratio, providing an up-to-2× tensor-math throughput path over dense matrix multiplication when the model is fine-tuned to respect the 2:4 constraint.

To understand why “Structured” patterns are required for hardware speedup, consider how sparse matrices are actually stored in memory. CSR and block sparse storage are treated formally in section C.1.5. Figure 11.7 compares their occupancy patterns and makes the block-index overhead visible. If the sparsity is random, the index overhead and irregular access kill performance. Structured sparsity, whether at the *large block* scale or the *fine-grained N:M* scale, makes this indexing predictable and compact, allowing hardware to fetch data efficiently.




Figure 11.7: Sparse Matrix Layouts: The four panels compare dense, CSR, block-sparse, and BSR matrix occupancy patterns. The separate nonzero-block-index column represents the metadata required to locate retained blocks in the block-sparse cases. Regular block structure keeps this metadata predictable and compact so hardware can skip zero blocks and fetch retained blocks efficiently.

The 2:4 pattern illustrates a broader principle: hardware achieves efficiency not by computing zeros faster, but by never loading them in the first place. This insight connects sparsity to the memory wall, since structured patterns reduce memory traffic, which is where the real cost lies.

Beyond structured sparsity optimizations, different hardware architectures implement matrix operations through distinct computational structures. Systolic arrays represent one such approach that has proven particularly effective for AI workloads.

11.4.6 Systolic arrays

While Tensor Cores package matrix operations into small, specialized instruction tiles, systolic arrays structure processing elements into large two-dimensional grids optimized for continuous data flow and operand reuse. By pulsing inputs horizontally and partial sums vertically through the array, systolic architectures minimize expensive DRAM fetches, enabling low-precision, quantized weights to achieve maximum compute density. The core motivation for systolic architectures stems from the same energy constraint that drives accelerator design: minimizing memory access penalties through physical operand reuse. A simple energy comparison through the array reveals why this architecture has become central to modern AI accelerators.

Napkin Math 11.1: The energy advantage of pulsing data

Scenario: The systolic architecture improves energy efficiency by keeping operands local as work pulses through the array; the “Systolic” (heartbeat) metaphor reflects this benefit of reusing operands locally. This worked model compares systolic dataflow with a naive implementation that streams every operand through DRAM:

1. **DRAM-streaming baseline:** Loads **A**, loads **B**, computes $\mathbf{A} \times \mathbf{B} + \mathbf{C}$, and writes **C** without cache or register reuse.
 - **Data movement:** 3 loads + 1 write = 4 DRAM accesses (per operation).
 - **Energy:** $\approx 4 \times 640 \text{ pJ} + 1 \text{ pJ (compute)} = 2561 \text{ pJ/op.}$
2. **Systolic Array (128 × 128 size):** Loads **A** and **B** once at the edges. Data “pulses” through 128 processing elements.
 - **Data movement:** 2 loads per 128 operations = 0.016 DRAM accesses (per operation).
 - **Energy:** $\approx 0.016 \times 640 \text{ pJ} + 1 \text{ pJ (compute)} \approx 11 \text{ pJ/op.}$

Systems insight: Under this deliberately naive DRAM-streaming baseline, the worked model gives the systolic array a 232.8× energy advantage. The ratio measures locality, not an inherent gap between vector and systolic arithmetic.

- Concretely, a 128×128 array can sustain 16,384 MACs/cycle with a large energy dividend by pulsing data through processing elements instead of repeatedly loading it from DRAM (Horowitz 2014; Jouppi et al. 2023).
- This locality lets dense arrays sustain many simultaneous MAC operations without paying a DRAM access for every operation.
- **Limitation:** The comparison assumes no effective reuse in the baseline and full reuse across the array. Caches, register blocking, array underutilization, and other dataflows change the ratio substantially.

A systolic array arranges processing elements in a grid pattern, where data flows rhythmically between neighboring units in a synchronized manner, enabling each operand to participate in multiple computations as it propagates through the array. This structured movement minimizes external memory accesses by maximizing local data reuse. A single weight value can contribute to dozens of operations as it moves through the processing elements, transforming the energy profile from memory-bound to compute-efficient execution.


Kung and Leiserson²⁶ (Kung and Leiserson 1979) first introduced systolic arrays, formalizing their use in parallel computing architectures for efficient matrix operations (Kung 1982). Unlike general-purpose execution units, systolic arrays exploit spatial and temporal locality by reusing operands as they propagate through the grid. Google’s TPU exemplifies this architectural approach: in the TPUv4, a 128×128 systolic array of multiply-accumulate units processes matrix operations by streaming data through the array in a pipelined manner (Jouppi et al. 2023). Figure 11.8 follows these data paths: a control unit feeds input buffers that stream data horizontally into the array, while the partial sums each cell produces flow vertically down to the accumulator chain at the bottom, which collects the finished results. Each processing element performs one multiply-accumulate per cycle and passes its operands to its neighbors, so a value loaded once is reused across an entire row or column rather than refetched from memory.

11.4.7 The tiling principle: Bridging graph and silicon

A fundamental mismatch exists between the *computational graph* (which sees a single $4,096 \times 4,096$ matrix multiplication) and the *physical silicon* (which possesses a fixed 128×128 systolic array). Bridging this gap requires tiling: the process of partitioning large tensor operations into “tiles” that fit exactly into the hardware’s fast local memory (SRAM or Scratchpad).

To process our 4,096-wide worked-example layer on a 128-wide systolic array, the compiler decomposes the result into 1,024 output tiles. Each output tile accumulates across 32 reduction-dimension tiles, so the complete matrix multiplication executes 32,768 **A**-tile-by-**B**-tile products. This is not merely a software convenience; it is a physical requirement. Operand tiles are fetched from slow HBM, staged in fast SRAM, and pulsed through the systolic array. Algorithm 11.1 states the loop nest a compiler emits for this decomposition: stream tiles of **A** and **B** on chip and accumulate their product into a tile of **C** before writing it back.

²⁶ | **Systolic array:** From Greek *sustole* (“contraction”), borrowed from cardiology to evoke rhythmic pumping. Kung and Leiserson used the term for arrays in which data advances through neighboring processing elements on a regular schedule (Kung and Leiserson 1979). Pipeline fill permits concurrent wavefronts and local reuse. The benefit depends on the stationary operand, tile, buffers, and utilization; it does not guarantee fixed DRAM savings. Irregular workloads and poorly aligned shapes remain a weaker fit because fill, drain, and idle elements consume capacity.



One extra dimension past the tile width tips utilization off a cliff.




Figure 11.8: Systolic Array Dataflow Architecture: A pipelined grid of processing elements (PEs) passes operands and partial results between neighbors. The illustrated movement is one possible matrix-multiplication dataflow; designs choose which operand or partial sum remains stationary to match reuse, buffers, and tensor shape. Utilization depends on mapping and dimensions.

The tile sizes are the lever. A larger tile reuses each loaded byte across more multiply-accumulate operations, raising the kernel’s arithmetic intensity and pushing it toward the compute-bound side of the roofline; the ceiling is how much of **A**, **B**, and **C** fits in fast on-chip memory at once. This tiling pattern is the central mechanism behind high-performance ML systems. It allows the hardware to maintain high system efficiency (η_{hw}) by ensuring that for every byte loaded from main memory, the data is reused 128× within the systolic grid. An engineer who understands tiling understands the “silicon contract”: if a layer’s dimensions are not multiples of the tile size (for example, a width of 129 on a 128 array), the system pays a *fringe tax* in underutilized silicon, where 127 units sit idle while one unit finishes the “remainder” tile.

Algorithm 11.1 Tiled (blocked) matrix multiply

Require: $\mathbf{A} \in \mathbb{R}^{M \times K}$, $\mathbf{B} \in \mathbb{R}^{K \times N}$; tile sizes T_M, T_N, T_K

Ensure: $\mathbf{C} = \mathbf{AB}$

```

1: for each row tile  $i_0$  (size  $T_M$ ) do
2:   for each column tile  $j_0$  (size  $T_N$ ) do
3:     initialize the C-tile accumulator to zero on chip
4:     for each  $k_0$  over  $K$  in steps of  $T_K$  do
5:       load A-tile  $[i_0, k_0]$  and B-tile  $[k_0, j_0]$  on chip
6:       accumulate the tile product on chip  $\triangleright$  reuse  $T_N/T_M$  per byte
7:     end for
8:     write the C-tile back to memory
9:   end for
10: end for

```

The systolic array architecture achieves computational efficiency through synchronized data movement across a structured grid of processing elements. Systolic arrays organize computation around four components:

- **Control unit:** Coordinates timing and data distribution across the array, maintaining synchronized operation throughout the computational grid.

- **Data streams:** Input matrices propagate through coordinated pathways where matrix A elements traverse horizontally while matrix B elements flow vertically through the processing grid.
- **Processing element grid:** Individual processing elements execute multiply-accumulate operations on streaming data, generating partial results that accumulate toward the final computation.
- **Output collection:** Results aggregate at designated output boundaries where accumulated partial sums form complete matrix elements.

The synchronized data flow ensures that matrix element A_{ik} encounters corresponding B_{kj} elements at precise temporal intervals, executing the multiply-accumulate operations required for matrix multiplication $C_{ij} = \sum_k A_{ik} \times B_{kj}$. This systematic reuse of operands across multiple processing elements substantially reduces memory bandwidth requirements by eliminating redundant data fetches from external memory subsystems.

For a 2×2 product, each output $C_{ij} = A_{i0}B_{0j} + A_{i1}B_{1j}$ requires two products and one accumulation. A weight-stationary mapping holds selected elements of **B** in local PE storage, streams elements of **A** across the array, and routes or retains partial sums until both k terms arrive. Thus C_{00} combines $A_{00}B_{00}$ with $A_{01}B_{10}$, while C_{01} combines $A_{00}B_{01}$ with $A_{01}B_{11}$. Values are injected with a time skew so operands belonging to the same product meet at the intended PE. The exact cycle path depends on array geometry, operand placement, and whether weights, inputs, or outputs are stationary. The lesson is local reuse: a loaded tile should serve several MACs before replacement, while fill and drain cycles remain part of the cost.

Each processing element in a 2D array performs a multiply-accumulate operation in every cycle. In a representative weight-stationary configuration:

1. Holds a stationary weight parameter ($B_{k,j}$) within its local register
2. Receives a streaming input activation ($A_{i,k}$) from its left neighbor
3. Computes the partial product $A_{i,k} \times B_{k,j}$ and adds it to the local accumulator
4. Passes the activation value rightward to the adjacent processing element for the next cycle

This structured computation model minimizes data movement between global memory and processing elements, improving both efficiency and scalability. 2D matrix multiplication operations map naturally onto grid-structured systolic processing arrays: follow the synchronized dataflow in figure 11.8, tracing how input activations stream horizontally while stationary weights accumulate partial products vertically. As systolic arrays operate in a streaming fashion, they are particularly effective for high-throughput workloads such as deep learning training and inference. However, as detailed in section 11.5.1, practical effectiveness is ultimately constrained by memory bandwidth bottlenecks.

A 128×128 systolic array capable of 16,384 operations per cycle requires a continuous operand stream to maintain utilization. On-chip buffers feed activations and weights to the array edges and are replenished from off-chip memory as needed; reuse within the array avoids fetching every operand from HBM each cycle. The TPU v4's 1,200 GB/s HBM2 bandwidth helps sustain this stream, but off-chip bandwidth can still limit models whose working sets and reuse patterns exceed on-chip capacity.

The quantization techniques in section 10.4 reduce model memory footprint by converting FP32 weights to INT8 representations. This optimization directly addresses the memory bandwidth constraints identified here. Converting 32-bit floating-point weights to 8-bit integers can reduce weight traffic by $4\times$; whether that changes a kernel from bandwidth bound to compute bound depends on the operation's original arithmetic intensity, the accelerator's INT8 ridge point (the intensity threshold at which its INT8 compute saturates), and the overhead of quantization and dequantization. Similarly, structured pruning removes entire rows or columns of weight matrices, reducing both the data volume that must traverse memory hierarchies and the computation required. These algorithmic optimizations prove valuable precisely because they target the memory bottleneck that limits accelerator performance in practice.

Systems Perspective 11.1: Matching architecture to workload

Challenge: Systolic arrays must choose which data to keep stationary (in registers) to minimize movement. This choice hard-codes the hardware’s preference for certain model types. Table 11.5 previews the three stationary-operand strategies and the workloads each favors; section 11.8 develops each one in full as a general mapping decision.

Table 11.5: Systolic-Array Dataflow Strategies: Three stationary-operand choices for systolic arrays, the reuse pattern each maximizes, and the workload class that benefits, showing how a fixed dataflow choice hard-codes an accelerator’s affinity for specific models.

Strategy	Stationary item	Optimized for	Example Workload
Weight-stationary	Weights (W)	High Reuse of Weights	CNNs (Conv2D): Filters are small and reused across the entire image.
Output-stationary	Partial Sums (C)	High Reuse of Accumulators	Large Batch MatMul: Accumulating results for many inputs against a large weight matrix.
Input-stationary	Inputs (A)	High Reuse of Activations	Transformers: The same activations feed many weight matrices across attention heads.

There is no “perfect” accelerator. A chip optimized for Weight-Stationary flow (like early TPUs) excels at CNNs where filters are small and heavily reused, but faces challenges with large language model (LLM) inference at small batch sizes, where the weight matrix is read once per token with minimal reuse, pushing architectures toward output-stationary or hybrid dataflow patterns.

11.4.8 Numerics in AI acceleration

Systolic arrays and Tensor Cores often support reduced-precision arithmetic. FP16 halves the bytes per value relative to FP32, and a target may provision more low-precision multiply-accumulate capacity. A $2\times$ speedup is not automatic: it requires a supported kernel and sufficient parallel work without another bottleneck. Building on chapter 10, reduced precision is a hardware-software decision. Input and accumulation formats may differ, affecting accuracy, throughput, energy, and movement.

11.4.8.1 Precision trade-offs

Lower precision is not free. FP16 has a 5-bit exponent and 10 stored fraction bits, giving it more significand precision than BF16 but much less dynamic range than FP32. Its smallest normal value is about 6.1×10^{-5} and its largest finite value is 65,504; subnormals extend lower with reduced precision. BF16 retains FP32’s 8-bit exponent and uses 7 fraction bits, preserving similar dynamic range at lower precision. INT8 has no floating-point exponent and depends on scale, zero-point, and clipping choices. Hardware and software must balance these properties against throughput and bandwidth.

The evolution of AI hardware reflects co-design between numerical methods and hardware capability. Earlier GPU generations lacked the dedicated mixed-precision matrix paths now used for deep learning, even when other units supported several formats. As training and inference methods demonstrated acceptable accuracy with selected lower-precision operations, vendors added native FP16, BF16, and integer tensor paths. Frameworks and libraries expose them through compatible kernels, autocasting, quantization, and calibration. Software gains materialize only when the model selects an efficient path; changing storage alone may reduce bytes without producing the advertised arithmetic speedup.

Precision support is integrated into execution units and memory paths. SIMD and SIMT lanes may support several scalar formats, while Tensor Cores (section 11.4.3) and systolic units (section 11.4.6) expose architecture-specific input and accumulation combinations. Lower precision reduces bytes per operand and may increase operations per cycle. Tiling and dataflow, not precision itself, determine reuse, while conversion, scaling, or fallback overhead can offset part of the gain. A format is useful only when the kernel and numerical policy can employ it safely.

Despite the advantages of reduced precision, deep learning models cannot always rely solely on low-bit representations. To address this challenge, modern AI accelerators implement mixed-precision computing, where different numerical formats are used at different stages of execution. These precision choices affect numerical reliability: matrix multiplications may be performed in FP16 or BF16, while accumulations are maintained in FP32 to prevent precision loss. Similarly, inference engines use INT8 arithmetic while preserving key activations in higher precision when necessary.

11.4.8.2 Mixed-precision computing

Modern AI accelerators increasingly support mixed-precision execution, allowing different numerical formats to be used at various stages of computation. Training workloads often use FP16 or BF16 for matrix multiplications, while maintaining FP32 accumulations to preserve precision (Miciekevicius et al. 2017; Mellempudi et al. 2019). The software implementation of mixed-precision training, including loss scaling techniques and framework support, is covered in section 8.6.3. Inference workloads, by contrast, optimize for INT8 or even INT4, achieving high efficiency while retaining acceptable accuracy.

The shift toward precision diversity is evident in the evolution of AI hardware. Early architectures such as NVIDIA Volta provided limited support for lower precision beyond FP16, whereas later architectures, including Turing and Ampere, expanded the range of supported formats. Table 11.6 traces this progression: Ampere GPUs introduced TF32 as a hybrid between FP32 and FP16 (NVIDIA 2020b), alongside broader support for BF16, INT8, and INT4 (NVIDIA Corporation 2017, 2018, 2020a).

Table 11.6: Precision Support Evolution: From Volta through Ampere, Tensor Core support expanded from FP16 to TF32, BF16, and integer formats, broadening the precision choices available to training and inference; the CUDA-core column shows the separate scalar/vector formats each generation supports.

Architecture	Year	Supported Tensor Core precisions	Supported CUDA Core Precisions
Volta	2017	FP16	FP64, FP32, FP16
Turing	2018	FP16, INT8, INT4, INT1	FP64, FP32, FP16, INT8
Ampere	2020	FP64, TF32, BF16, FP16, INT8, INT4	FP64, FP32, FP16, BF16, INT8

Newer architectures incorporate a growing diversity of numerical formats because different workloads bind at different points on the accuracy-throughput-energy trade-off. Precision support is therefore another form of workload matching, not a generic feature checklist.

The precision format used in hardware design has cascading implications across the entire system. Reducing from FP32 to FP16 cuts memory traffic in half, which matters far more than it might seem: because memory access dominates energy consumption, halving memory traffic can substantially reduce energy per inference when data movement is the bottleneck (Horowitz 2014). Simultaneously, Tensor Cores and systolic arrays can pack more lower-precision multiply-accumulate units into the same silicon area, raising peak throughput (Dally et al. 2021; Dally 2023). Lower-precision integer arithmetic also consumes less energy than FP32 in representative technology studies (Horowitz 2014), and the inference-focused TPUv1 was built around 8-bit multiply units (Jouppi et al. 2017). The systems insight is that reduced precision does not merely “save bits”: it simultaneously relieves the memory bandwidth bottleneck *and* increases compute density, attacking both sides of the roofline at once.

As AI models continue to scale, precision support connects the compute primitive discussion back to the memory wall: lower-bit formats matter when they reduce the bytes moved and keep the hardware’s matrix engines fed. The remaining architectural question is how these execution units, precision formats, and memory paths integrate into complete accelerator systems. Architectural integration determines how efficiently computational primitives become usable accelerator throughput. SIMD lanes, Tensor Cores, and systolic arrays are building blocks, but their full-chip organization varies significantly across AI processors; the choice of execution units, their numerical precision support, and their connectivity shape how effectively hardware can scale for deep learning workloads.

11.4.9 Intra-node interconnects: Scaling the stack

Mastery of the single-machine stack requires understanding how bits move between GPUs and the CPU. In the 1–8 GPU regime, scaling is achieved through high-speed intra-node interconnects such as NVLink and host-to-device PCIe transfers that mitigate the memory wall. These links form a **bandwidth taper**: data-movement speed falls at each step away from the compute units, from on-package HBM through the GPU-to-GPU NVLink bridge down to the host PCIe link. The PCIe step is much slower than the accelerator-local memory and inter-GPU fabric, so any data path that touches the CPU can become a performance hazard, the “PCIe Wall” that NVLink exists to avoid (NVIDIA Corporation 2020c). Section 11.5.5.1 develops this hierarchy quantitatively, where host-accelerator communication is the operative concern.

Modern AI processors exhibit a range of design trade-offs based on their intended applications, and comparing their configurations reveals how deployment constraints drive architectural divergence. A training-optimized accelerator like the NVIDIA A100 packs many Streaming Multiprocessors with wide SIMD units and FP16 Tensor Cores because training throughput scales with aggregate multiply-accumulate capacity (NVIDIA Corporation 2020a). Google’s TPUv4 makes a radically different bet: just two cores per chip, each containing massive BF16 systolic arrays, a design that trades programmer flexibility for efficiency on dense matrix multiplications (Jouppi et al. 2023). At the inference end, Intel’s Sapphire Rapids dedicates Advanced Matrix Extensions (AMX) tile engines to INT8 and BF16, reflecting the insight from chapter 10 that inference models tolerate reduced precision (Intel Corporation 2021a). Mobile neural engines take this further by shrinking matrix engines into low-power system-on-chip (SoC) blocks, prioritizing energy efficiency per operation over peak throughput. Table 11.7 compares these architectural configurations.

Table 11.7: AI Processor Configurations: Modern AI processors prioritize different execution unit characteristics for specific workloads: NVIDIA A100 uses wide SIMD and Tensor Cores for training, Google TPUv4 emphasizes high-throughput BF16 matrix multiplication, Intel Sapphire Rapids focuses on INT8-optimized inference, and mobile NPUs prioritize low-power execution. These variations in SIMD width, Tensor Core size, and processing element count reflect the growing diversity in AI hardware architectures.

Processor	Vector execution	Matrix engine	Processing Elements	Primary workloads
NVIDIA A100	CUDA FP/INT lanes	Tensor Core instruction tiles	108 SMs	Training, HPC
Google TPUv4	Vector unit	128×128 BF16 systolic arrays	2 cores/chip	Training
Intel Sapphire	512-bit AVX-512	AMX tiles (up to 16 rows by 64 bytes each)	Up to 60 cores	Inference
Mobile NPU	CPU/GPU/DSP vectors	Small matrix engines	Integrated NPU blocks	Mobile inference

The pattern across these configurations reveals a consistent engineering principle: each design sacrifices generality to optimize for its target workload’s dominant operation and precision. Training chips invest silicon in wide floating-point datapaths; inference chips trade precision for throughput; mobile chips trade throughput for energy efficiency. No single design dominates across all workloads, which is precisely why hardware selection depends on workload analysis rather than headline specifications.

11.4.10 Cost-performance analysis

Architectural specifications define computational potential, but deployment decisions require cost-performance analysis. Raw compute is only one input: delivered performance may be limited by arithmetic, memory, communication, software support, or utilization.

The energy differential established in section 11.2.5 (where memory access costs dominate computation) drives the entire specialized hardware revolution. This disparity helps explain why many accelerators achieve only a fraction of peak compute on memory-bound workloads, while architectures that maximize data reuse (for example, systolic arrays on dense matrix kernels) can sustain substantially higher utilization under favorable conditions.

Consider an organization choosing between more older accelerators and fewer newer ones. Peak FLOP/s can mislead whenever important kernels are bandwidth or communication bound. Large

training matrix multiplications may be compute bound, while optimizer steps, embeddings, normalization, and some inference paths may not be. The relevant comparison is delivered workload throughput per dollar under the expected utilization, memory capacity, bandwidth, and scaling efficiency.

These dynamics help explain the rapid adoption of newer accelerators despite higher unit prices. For memory-bound workloads, improvements in effective bandwidth (and the software stack's ability to use it) can dominate real-world performance. Cloud deployment further complicates the analysis, as rental pricing, utilization, and operational overheads can change the break-even point between purchasing and renting hardware.

Table 11.8 provides an illustrative cost worksheet for common accelerators. Prices vary by vendor, region, contract, and purchase volume. More importantly, the throughput rows use different precision modes, so price per peak operation is meaningful within a row but not a controlled cross-device benchmark. The table is a prompt to make assumptions explicit, not a ranking.

Table 11.8: Illustrative Accelerator Cost Worksheet: Each row divides an approximate acquisition-price assumption by the representative peak mode shown. Because precision, sparsity assumptions, and price sources differ, the ratios are not comparable as workload performance. Total cost also includes utilization, memory capacity, power, cooling, interconnect, software, and infrastructure; TPU acquisition cost is estimated from cloud pricing.

Accelerator	List price	Representative Peak Throughput (precision shown)	Memory Bandwidth	Price/performance
NVIDIA V100	~\$10,000	125 TFLOP/s	900 GB/s	\$80/(TFLOP/s)
NVIDIA A100	~\$15,000	312 TFLOP/s	2,039 GB/s	\$48.1/(TFLOP/s)
NVIDIA H100	~\$25,000–30,000	494 TFLOP/s (TF32)	3,350 GB/s	~\$50.6/(TFLOP/s)
Google TPUv4	~\$8,000*	275 TFLOP/s (BF16)	1,200 GB/s	~\$29.1/(TFLOP/s)
Intel Gaudi 2	~\$12,000	865 TFLOP/s (FP8)	2,450 GB/s	\$13.9/(TFLOP/s)

The worksheet shows why a hardware decision cannot be reduced to one column. A bandwidth-bound kernel may value H100's 3,350 GB/s memory path more than another increase in peak arithmetic, while a compute-bound kernel may show the reverse. Capacity can determine whether a model fits at all, and software support can determine whether the advertised path is reachable. The next section isolates the memory side of that decision.

Framework selection significantly impacts these economic decisions. Detailed hardware-framework optimization strategies are covered in chapter 7, while performance evaluation methodologies are discussed in chapter 12.

The preceding sections introduced vector instructions, matrix-tile operations, and Tensor Cores capable of very high peak throughput. An NVIDIA A100 advertises 312 TFLOP/s on its FP16 tensor path, and newer accelerators add FP8 paths (NVIDIA Corporation 2020a; Kuzmin et al. 2022; Micikevicius et al. 2022). Dividing a ResNet-50 operation count by that peak yields a microsecond-scale arithmetic lower bound, not an inference-time prediction.

NVIDIA's Blackwell (B200) architecture extends this trend with an FP4 tensor path, advertised at up to 9 PFLOP/s dense or 18 PFLOP/s sparse peak throughput per chip (NVIDIA Corporation 2024). Reaching those peaks requires a model, numerical method, kernel, and sparsity pattern that use the corresponding path. Lower-bit hardware expands the design space; it does not make every workload safely reducible to FP4.

Yet real ResNet-50 inference takes milliseconds, not microseconds. The gap between theoretical capability and practical performance reveals the chapter's central tension, first posed in the Purpose section: computational capability has outpaced our ability to feed data to processors. *Moving data from memory costs orders of magnitude more energy than arithmetic, and memory bandwidth has improved more slowly than tensor arithmetic throughput.* This disparity determines whether those 312 TFLOP/s translate into low sustained utilization or high sustained utilization on a particular workload (Horowitz 2014; Gholami et al. 2024).

Understanding this gap requires examining the memory systems that feed the compute primitives. The memory hierarchy is not merely supporting infrastructure; together with arithmetic intensity and mapping, it determines how much peak throughput a workload can use.

11.5 AI Memory Systems

ResNet-50 can expose the gap between accelerator arithmetic and memory: convolution weights, activations, and intermediate results still have to arrive on time. Modern accelerators reach hundreds of TFLOP/s or more on selected low-precision paths (NVIDIA Corporation 2020a, 2024; Choquette 2023), but a kernel cannot use that capacity when its required traffic exceeds the relevant memory bandwidth. The **AI memory wall** names this growing mismatch between arithmetic demand and the memory system that feeds it.

Unlike conventional workloads, ML models require frequent access to large volumes of parameters, activations, and intermediate results, leading to substantial memory bandwidth demands. This challenge intersects with the data management strategies covered in chapter 4. Modern AI hardware addresses these demands through advanced memory hierarchies, efficient data movement techniques, and compression strategies that promote efficient execution.

11.5.1 Understanding the AI memory wall

Definition 11.4: AI memory wall

*The **AI memory wall** is the ML accelerator performance constraint that arises when arithmetic throughput (R_{peak}) outpaces memory bandwidth (BW). The core issue is whether the memory system can supply operands and absorb results quickly enough to sustain the specialized primitives introduced in section 11.4.*

1. **Significance:** A workload has reached the memory wall when increasing FLOP/s alone no longer improves it because the $\frac{D_{\text{vol}}}{\text{BW}}$ term dominates the relevant execution time.
2. **Distinction:** Unlike a general-purpose memory wall, which affects all computing, the AI memory wall is driven by the massive model state and activation storage required by deep learning.
3. **Common pitfall:** A frequent misconception is that the memory wall is “fixed” by more memory. In reality, it is a bandwidth-latency gap: even with infinite capacity, the speed of moving data between memory and compute remains the fundamental physical bottleneck.

Data access can dominate the energy budget of a memory-intensive kernel.²⁷ Figure 11.9 compares representative operation costs from one technology study, revealing the multi-order-of-magnitude gap between simple local arithmetic and an off-chip DRAM access.

11.5.1.1 Quantifying the compute-memory performance gap

The energy disparity that figure 11.9 captures grows more severe with each hardware generation. Over the past two decades, peak computational capabilities have grown substantially faster than DRAM bandwidth (Gholami et al. 2024). This divergence creates a widening gap where accelerators possess massive computational power but cannot access data quickly enough to use it. Representative high-end accelerators can deliver on the order of 10^3 TFLOP/s of peak tensor throughput (for example, NVIDIA H100 delivering 989 TFLOP/s in FP16 or nearly 2,000 TFLOP/s in dense FP8) while providing approximately 3.35 TB/s of memory bandwidth (Choquette 2023). This implies that on the order of 10^2 FLOP of work per byte moved is required to fully use the compute, which can exceed the arithmetic intensity of many practical neural network workloads.

The memory wall manifests through three critical constraints. First, the energy disparity: accessing DRAM can consume orders of magnitude more energy than a multiply-accumulate operation (Horowitz 2014; Sze et al. 2017), which often shifts bottlenecks from raw compute to power and data movement. Second, the bandwidth limitation: even TB/s memory systems may not feed large parallel compute arrays continuously on memory-bound workloads, leaving compute underutilized. Third, the latency hierarchy: off-chip memory access can require hundreds of cycles, creating pipeline stalls that cascade through parallel execution units.

²⁷ | **Von Neumann bottleneck:** Separating storage from compute requires instructions and data to cross interfaces rather than remain in local registers. In the representative 45 nm estimates from Horowitz (2014), a DRAM access consumes over $20,000\times$ the energy of an INT8 addition. The exact ratio changes with technology and operation, but locality remains a first-order design concern.




Figure 11.9: The Energy Hierarchy: Representative 45 nm energy estimates on a logarithmic scale from Horowitz (2014). An off-chip DRAM access costs about $128\times$ an 8 KB SRAM read and over $20,000\times$ an INT8 addition in this model. The values are not universal device specifications; they illustrate why reuse in registers and SRAM can matter more than reducing arithmetic count alone.

11.5.2 Hardware balance (I_{ridge}): The paradigm partition

Different devices place the compute-bandwidth boundary at different intensities. The **hardware balance** (I_{ridge}) is the ratio of peak arithmetic throughput to peak memory bandwidth. In the roofline model, it is the ridge point where the compute and bandwidth ceilings intersect:

$$I_{\text{ridge}} = \frac{R_{\text{peak}}}{\text{BW}}$$

The ratio sets a device-specific boundary. Depending on the chosen precision, a high-end accelerator such as H100 may require roughly 150–300 FLOP/byte to reach its arithmetic ceiling, whereas a microcontroller may have a much lower ridge point. The same kernel can therefore be bandwidth bound on a compute-dense accelerator and compute bound on an edge processor. Hardware balance does not label either device as universally efficient; it identifies how much reuse a workload needs on that device.

Peak computational capability has expanded faster than memory bus performance across accelerator generations, and the size of that mismatch is what sets how much data reuse a kernel needs to stay fed. Figure 11.10 separates the two growth rates that produce the ridge point, each normalized to the V100 baseline, so the widening distance between them can be read directly.

The threshold required for a kernel to become compute bound has shifted across accelerator hardware generations. Follow the upward trajectory of hardware ridge points in figure 11.11 from V100 through B200, observing how the arithmetic intensity requirement increases over time.

Beyond performance limitations, memory access imposes a steep energy cost. Fetching data from off-chip DRAM consumes far more energy than performing arithmetic operations (Horowitz 2014). This inefficiency is particularly evident in machine learning models, where large parameter sizes, frequent memory accesses, and nonuniform data movement patterns exacerbate memory bottlenecks. The energy differential drives architectural decisions: Google’s TPUv1 achieved $30\text{--}80\times$ better performance per watt than contemporary CPUs and GPUs on Google’s inference benchmarks by minimizing data movement through systolic arrays and large on-chip memory (Jouppi et al. 2017). These design choices demonstrate that energy constraints, not computational limits, often determine practical deployment feasibility. Tracing a single tensor through every level of the memory hierarchy during a real inference pass makes these energy costs concrete.

11.5.2.1 Memory access patterns in ML workloads

Beyond raw computational throughput, an accelerator’s efficiency depends on its ability to continuously supply data to processing units without stalls. Neural networks impose three concurrent




Figure 11.10: The Compute-Bandwidth Divergence: Peak dense FP16 throughput and memory bandwidth across five NVIDIA datacenter generations, each normalized to its V100 baseline on a log scale. Over 2017–2024 compute grows 18× while bandwidth grows 8.9×, so the shaded separation between the curves widens by 2.0×. That separation is the data reuse a kernel must supply to keep the arithmetic units busy. The H200 step is the instructive exception: bandwidth rose while peak compute held flat, narrowing the gap for one generation.




Figure 11.11: The Rising Ridge: Hardware arithmetic intensity (FLOP/byte) over time computed as each chip’s dense FP16 tensor peak divided by its memory bandwidth. This is the arithmetic intensity a kernel must reach before compute rather than bandwidth limits it. As compute capability grows faster than memory bandwidth, the ridge point rises from V100 through H100 and remains high on B200. This trend explains why architectures with high data reuse flourish while low-reuse workloads face a growing hardware tax.

demands on this data supply. Model parameters (weights and biases) may number in the billions, requiring efficient storage and streaming to maintain throughput. Intermediate activations produced at each layer must be temporarily held for subsequent operations, contributing to memory overhead in deep architectures. During training, backpropagation adds a third demand: storing and accessing gradients for every parameter, further increasing data movement volume between compute units and memory.

As models increase in size and complexity, improvements in memory capacity and bandwidth become increasingly important. Although specialized compute units accelerate operations like matrix multiplications, their overall performance depends on the continuous, efficient delivery of data to the processing elements. In large-scale applications such as natural language processing and computer vision, models often incorporate millions to billions of parameters (Brown et al. 2020),

and achieving high performance requires minimizing delays and stalls caused by inefficient data movement between memory and compute units (Narayanan et al. 2021; Kwon and Rhu 2018).



Lighthouse 11.2: Life of a tensor: GPU-hosted KWS

Recall the Keyword Spotting lighthouse summarized in table 2.2. If the same one-second audio clip is batched on a GPU-hosted inference path, its physical journey through the memory hierarchy looks like this:

1. **DRAM (HBM):** The tensor starts here.
 - **Size:** 16,000 samples $\times$ 2 bytes (FP16) = 32 KB.
 - **Latency:** Fetching this from off-chip memory takes ~ 300 ns (plus queuing delay).
 - **Energy:** Cost is ~ 20 pJ/bit. High cost.
2. **L2 cache:** Device-memory requests may be served here when the requested cache lines are resident.
 - **Latency:** ~ 4 ns.
 - **Access:** Shared across multiple Streaming Multiprocessors (SMs).
3. **L1 cache/shared memory:** A specific SM works on a tile of the audio, using hardware-managed cache or explicitly staged shared memory.
 - **Latency:** ~ 1 ns.
 - **Locality:** If an L1 request misses, it can still be served by L2 before an HBM access is required.
4. **Registers:** The Tensor Core operates here.
 - **Latency:** Typically a small number of cycles, depending on the instruction and architecture.
 - **Throughput:** 312 TFLOP/s.
 - **Energy:** Cost is ~ 0.1 pJ/bit.

Systems insight: Performance is bounded by the lower of peak compute throughput and bandwidth times arithmetic intensity. Which memory link supplies the relevant bandwidth depends on where the working set resides; HBM-to-L2 is one possible limiting path, not the roofline's sole determinant.

One way to quantify this challenge is by comparing the data transfer time with the time required for computations. To do this, five key variables are defined: D_{vol} is the total data volume (bytes), BW is the available memory bandwidth (bytes/s), O is the number of floating-point operations, R_{peak} is the peak hardware throughput (FLOP/s), and η_{hw} is the realized hardware utilization.

The memory transfer time T_{mem} and compute time T_{compute} are expressed as:

$$T_{\text{mem}} = \frac{D_{\text{vol}}}{BW}$$

$$T_{\text{compute}} = \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}$$

In a no-overlap model, $T_{\text{mem}} > T_{\text{compute}}$ indicates memory-bound execution. With overlap, the larger term often dominates elapsed time, but dependencies, queueing, and incomplete overlap determine how much compute actually sits idle and which data-movement optimization can improve performance.

Figure 11.12 quantifies this disparity for specific public-count models and hardware generations. The gap between these curves represents the engineering challenge that drives accelerator memory-system design (Krizhevsky et al. 2012; Brown et al. 2020; Chowdhery et al. 2022; Dubey et al. 2024). Even high-bandwidth accelerators like NVIDIA's B200 and AMD's MI300X/MI325X-class devices

cannot close this growth-rate gap by bandwidth alone (NVIDIA Corporation 2024; AMD 2023). Parameter counts for proprietary systems such as GPT-4 and Gemini are not officially disclosed, so they are not plotted as factual data points.




Figure 11.12: Model Size vs. Hardware Bandwidth: Both series are the base-10 logarithm of growth relative to 2014, so each unit on the vertical axis is a tenfold increase. Over the plotted decade parameter counts peak at roughly 4,900× against bandwidth’s 17×. The shaded band is a gap between logarithms rather than a difference, so its width reads as a ratio, and the two fitted trends diverge at a constant rate.

11.5.2.2 Irregular memory access

Many ML workloads combine regular dense kernels with irregular memory pressure from sparsity, embedding lookups, variable sequence lengths, attention/KV-cache traffic, and small batches. The dense parts are exactly why accelerators work so well; the irregular parts are where standard caching mechanisms and memory hierarchies struggle, leading to increased memory latency and inefficient bandwidth utilization.

Many ML systems contain both access regimes. Table 11.9 contrasts regular dense kernels with irregular components such as embeddings, sparse operators, and input-dependent routing.

One key source of irregularity in ML workloads stems from batch size and execution order. The way input data is processed in batches directly affects memory reuse, creating a complex optimization challenge. Small batch sizes decrease the likelihood of reusing cached activations and weights, resulting in frequent memory fetches from slower, off-chip memory. Larger batch sizes can improve reuse and amortize memory access costs, but simultaneously place higher demands on available memory bandwidth, potentially creating congestion at different memory hierarchy levels. This delicate balance requires careful consideration of model architecture and available hardware resources.

Table 11.9: Regular and Irregular Memory Access: Regular dense kernels expose tileable reuse and predictable addresses, whereas irregular ML components rely more on indirect or input-dependent access, weakening caching and prefetching and adding metadata or load-balance overhead. A single model may contain both regimes.

Feature	Regular dense kernels	Irregular ML components
Access pattern	Contiguous or tiled	Indirect, scattered, or input-dependent
Cache locality	Often high after tiling	Often lower and workload-dependent
Data reuse	Structured across tiles or batches	Sparse or dynamic, depending on the operator
Dependencies	Static loop nests	Input-dependent routing or variable shapes
Examples	GEMM, convolution, dense attention	Embedding lookup, unstructured sparsity, MoE routing
Primary pressure	Compute or bandwidth, depending on arithmetic intensity	Latency, bandwidth, metadata, and load imbalance
Energy consequence	Depends on reuse and memory tier	Extra movement and metadata can increase energy

28 | Sparsity and memory irregularity: Unstructured sparse formats commonly gather nonzero elements through indirect addressing, which can weaken sequential access and make metadata overhead significant. Structured sparsity preserves hardware-friendly blocks, so its access behavior differs from arbitrary pruning. Without suitable kernels or hardware support, an unstructured sparse model can run slower than its dense counterpart despite performing fewer arithmetic operations.

Different neural network layers interact with memory in distinct ways beyond batch size considerations. Convolutional layers benefit from spatial locality, as neighboring pixels are processed together and small weight kernels are reused. Fully connected layers and dense attention use regular matrix accesses, but their large weights or key-value tensors can exceed cache capacity and become bandwidth bound when reuse is low. Variable sequence lengths complicate memory planning without making dense attention accesses intrinsically random.

Irregular access can instead arise from unstructured sparsity,²⁸ embeddings, and input-dependent routing. Structured sparsity preserves regular blocks that specialized hardware can process efficiently, whereas compressed unstructured formats often require indirect or scattered accesses. Mixture-of-Experts routing is deterministic for a given input but varies across inputs, complicating batching, prefetching, and load balance.

Irregularity degrades performance through distinct mechanisms, and treating them as one memory problem can point to the wrong optimization. Scattered addresses may turn one logical access into several memory transactions, so each fetched cache line carries few useful payload bytes. Indirection can also serialize address generation: the system must fetch an index or pointer before it knows which payload to request, limiting memory-level parallelism even when aggregate bandwidth remains available. Input-dependent routing creates a third failure mode when some processing elements receive longer queues than others; completion time follows the slowest queue while neighboring units idle. Variable-sized allocations may fragment capacity, but fragmentation is neither necessary nor sufficient for irregular access. A sequential stream can saturate bandwidth without being irregular, whereas a dependent lookup chain can stall far below peak bandwidth.

Diagnosis should therefore begin with the mechanism that prevents forward progress. A low ratio of useful payload bytes to transferred bytes indicates poor spatial utilization or coalescing. Long memory stalls alongside unused bandwidth indicate dependent misses or too few concurrent requests to hide latency. Uneven queue depths or per-route work reveal load imbalance rather than a cache-capacity problem. Cache hit rate alone cannot distinguish these cases because a miss can waste bandwidth, expose latency, or both. Each diagnosis implies a different remedy: packing or blocking improves spatial utilization, reordering and prefetching expose concurrent requests, structured sparse formats reduce indexing overhead, and routing or partitioning balances work. The memory hierarchy determines the cost of the remaining accesses, while dataflow and tensor-layout choices determine how much irregularity software can remove before execution reaches the hardware.

The optimization order follows that diagnosis. Representation comes first when the format is under system control because a block-structured sparse format or grouped routing can remove indirection before a kernel executes. Request organization comes next: accesses issued together should target adjacent addresses, and enough independent requests must remain in flight to hide unavoidable latency. Additional cache capacity or bandwidth helps only when profiling shows that the regularized request stream still exhausts that resource. More bandwidth cannot resolve a serial pointer dependency, and more capacity cannot rebalance skewed work. The objective is to expose as much regular structure as possible while confining unavoidable irregularity to the smallest portion of execution.

The irregular access patterns and memory wall constraints examined in section 11.5.1 create formidable challenges, but they also reveal optimization opportunities. Although individual memory accesses may appear unpredictable, ML workloads exhibit structured reuse patterns at a higher level: the same weights are applied across batch elements, the same kernels slide across spatial dimensions, and the same attention patterns recur across sequence positions. Hardware designers exploit these regularities through carefully structured memory hierarchies that maintain frequently accessed data close to compute units, even when the specific access sequence varies.

11.5.3 Memory hierarchy

Modern AI accelerators exploit these structured reuse patterns through multilevel memory hierarchies: rather than treating memory as a monolithic resource, they organize storage into distinct tiers optimized for different access patterns, reuse distances, and energy costs. While general-purpose computing contends with unpredictable memory access, ML workloads exhibit structured reuse that can be optimized through careful data organization across multiple memory levels.

At the highest level, large-capacity but slow storage devices provide long-term model storage. At the lowest level, high-speed registers and caches ensure that compute units can access operands with minimal latency. Between these extremes, scratchpad memory provides software-managed on-chip storage, while accelerators use device-local DRAM, such as HBM or GDDR, for larger working sets.

The key pattern in table 11.10 is that larger, more distant storage generally provides greater capacity at higher latency and energy cost. The ratios are not uniform between adjacent levels, and nearby register or SRAM accesses need not cost more than every arithmetic operation.

Table 11.10: Memory Hierarchy Trade-Offs: AI accelerators use a multilevel memory hierarchy to balance performance and capacity. Each level provides distinct latency, bandwidth, and capacity characteristics that dictate how neural network components (weights, activations, and intermediate results) should be allocated to minimize bottlenecks and maximize throughput. Relative labels describe the typical ordering; exact values depend on architecture. Section D.3.1 supplies registry-backed reference latencies.

Memory Level	Relative latency	Bandwidth	Relative capacity	Example Use in Deep Learning
Registers	Lowest	Highest	Tiny	Storing operands for immediate computation
L1/L2 Cache (SRAM)	Very low	High	Small	Caching frequently accessed activations and small weight blocks
Scratchpad memory	Very low	High	Small to medium	Software-managed storage for intermediate computations
Device DRAM (HBM, GDDR, or LPDDR)	High	Moderate–Very High	Large	Storing model parameters and activations that do not fit on-chip
Host DRAM (typically DDR)	High locally; transfer adds latency	Moderate	Large to very large	Staging or offloading data outside accelerator memory
Flash Storage (SSD/NVMe)	Very high	Low	Very large	Storing pretrained models and checkpoints for later loading

The hierarchy invites an apparently simple solution: build larger, faster off-chip memory and eliminate the need for on-chip SRAM entirely. The answer is rooted in physics: signal propagation within and between chips imposes a hard latency floor.

Napkin Math 11.2: The speed of light limit

Problem: Why is on-chip SRAM necessary instead of fetching all data from HBM?

Math:

1. **Distance:** On an H100-class 814 mm² die, signals travel ~20 mm.
2. **Propagation latency:** At $\approx 0.5c$ in silicon, $20 \text{ mm} / (0.5c) \approx 130 \text{ ps}$.
3. **Clock cycle:** At 2 GHz, $1 / (2 \text{ GHz}) = 500 \text{ ps}$.
4. **DRAM:** Off-chip HBM sits millimeters away on the package, but DRAM access latency plus protocol overhead = 100+ cycles.

Systems insight: The 20 mm propagation estimate is about 130 ps, less than one 500 ps clock cycle, so distance alone does not prove a one-cycle access impossible. A real HBM access includes DRAM-array latency, serialization, protocol traversal, interconnect delay, and queuing, producing this 100-plus-cycle access path. Local registers and SRAM are therefore required to feed compute units at 2 GHz. For transformer models, this complete access path is why weights must be staged in SRAM in tiles rather than read on demand from HBM, and why repeatedly streaming an entire KV cache can collapse inference throughput when its memory traffic cannot be hidden behind arithmetic.

11.5.3.1 On-chip memory

On-chip memory includes registers, Static Random-Access Memory (SRAM) caches, and software-managed scratchpads on the accelerator die. Each level trades capacity, latency, bandwidth, and

management policy. Aggregate on-chip bandwidth can reach many TB/s, but the value depends on access width and concurrency, and capacity is limited and partitioned across cores or tiles. Device DRAM such as HBM provides much greater capacity at lower bandwidth and higher latency; its delivered rate likewise depends on access pattern. Registers provide the most immediate operand and accumulator storage, but they are not limited to only a few operands across the chip. A finite register file is divided among active threads or instructions, so per-thread allocation affects occupancy and tile size. Efficient kernels retain short-lived values in registers while moving reusable tiles through caches or scratchpads, avoiding spills and lower-tier transfers when capacity permits.

To reduce the need for constant data movement between registers and external memory, small but fast caches serve as an intermediary buffer. These caches store recently accessed activations, weights, and intermediate values, ensuring that frequently used data remains available with minimal delay. However, the size of caches is limited, making them insufficient for storing full feature maps or large weight tensors in machine learning models. As a result, only the most frequently used portions of a model's parameters or activations can reside here at any given time.



Example 11.3: The Tensor Core contract

Scenario: An engineering team migrates a transformer workload to NVIDIA A100 GPUs expecting automatic Tensor Core acceleration, but profiling reveals Tensor Cores remain idle (NVIDIA Corporation 2020a).

Diagnosis: The workload used shapes and data types that selected less efficient kernels rather than the intended Tensor Core path.

Systems lesson: Specialized paths have operator, precision, layout, and shape requirements. Libraries can often handle nonaligned dimensions through padding or alternative kernels, but reaching high utilization requires verifying which path was actually selected.

For larger working datasets, many AI accelerators include scratchpad memory, which offers more storage than caches but with a key difference: it allows explicit software control over what data is stored and when it is evicted. Unlike caches, which rely on hardware-based eviction policies, scratchpad memory enables machine learning workloads to retain key values such as activations and filter weights for multiple layers of computation. This capability is useful in models like convolutional neural networks, where the same input feature maps and filter weights are reused across multiple operations. By keeping this data in scratchpad memory rather than reloading it from external memory, accelerators can significantly reduce unnecessary memory transfers and improve overall efficiency (Chen et al. 2017). On NVIDIA GPUs, the hardware exposes scratchpad memory to programmers as *shared memory*: a fast, software-managed SRAM region that all threads in a thread block can read and write, distinct from the hardware-managed L1/L2 caches. Custom ML kernels written in CUDA or Triton control this memory explicitly. FlashAttention achieves its substantial throughput gains for transformer attention layers precisely by exploiting this mechanism: rather than materializing the full $S \times S$ attention score matrix in HBM, it tiles queries, keys, and values through SRAM/shared memory and writes only the final output back to HBM (T. Dao et al. 2022). The reduction in HBM round-trips—not fewer arithmetic operations—is the primary source of the speedup.

11.5.3.2 Off-chip memory

Once a model's working set outgrows on-chip SRAM, the design question is whether it fits in the accelerator's device-local DRAM. An accelerator may use HBM, GDDR, or LPDDR; these are alternative memory technologies rather than sequential tiers. Data that does not fit must be staged or offloaded through host memory or storage, adding interconnect transfers before the accelerator can use it.

Beyond on-chip memory, HBM provides rapid access to larger model parameters and activations that do not fit within caches or scratchpad buffers. It stacks multiple memory dies and uses wide interfaces to deliver much higher aggregate bandwidth than conventional off-chip DRAM interfaces. It is often used for model parameters and activations on high-performance accelerators, while its

cost, packaging complexity, and power requirements make it less common in constrained edge devices.

Accelerators without HBM may instead use device-local GDDR or LPDDR. When a model exceeds device-memory capacity, systems can offload data to host DRAM, but the accelerator must move that data across an interconnect such as PCIe before computation. Effective offloading therefore stages only the portions needed for each computation and overlaps transfers when possible.

At the highest level of the hierarchy, flash storage and solid-state drives (SSDs) store large pre-trained models, datasets, and checkpointed weights. These storage devices offer large capacities but are too slow for real-time execution, requiring models to be loaded into faster memory tiers before computation begins. For instance, in training scenarios, checkpointed models stored in SSDs must be loaded into DRAM or HBM before resuming computation, as direct execution from SSDs would be too slow to maintain efficient accelerator utilization (Narayanan et al. 2021).

The memory hierarchy thus balances competing objectives of speed, capacity, and energy efficiency. Moving data through multiple levels introduces latency and bandwidth costs, while limited capacity forces additional movement when working sets exceed local storage. Memory bandwidth is therefore a major determinant for bandwidth-bound workloads, alongside compute throughput and software overhead.

11.5.4 Memory bandwidth and architectural trade-offs

Advertised memory bandwidth is only a ceiling; achievable bandwidth depends on access pattern, batching, locality, and the host interface that feeds the accelerator. Modern accelerators exhibit distinct bandwidth-capacity trade-offs that directly shape which workloads they can serve efficiently. Representative data center accelerators provide memory bandwidth on the order of a few TB/s, often paired with tens of GB of high-bandwidth memory. *Raw bandwidth* alone, however, is misleading: what matters is *achievable bandwidth* for a given access pattern. Transformer attention, convolution, and fully connected layers can all realize different fractions of peak bandwidth because their reuse, tiling, and access regularity differ. Fully connected layers approach peak bandwidth only when batch sizes are large enough to amortize the cost of loading weight matrices—which connects directly to the batch-size sensitivity discussed in the roofline analysis in section 11.6. The practical consequence is that an accelerator’s effective bandwidth for a specific workload may be well below its advertised peak, making bandwidth-per-dollar a more reliable purchasing metric than peak bandwidth alone.

As established in section 11.5.1, on-chip memory access typically consumes energy in the single-digit-to-tens of picojoules per access, while external DRAM can be on the order of hundreds of picojoules per access, an orders-of-magnitude energy penalty. AI accelerators minimize DRAM access through three key strategies: weight stationarity (keeping model parameters in on-chip memory), input stationarity (buffering input activations locally), and output stationarity (accumulating partial sums on-chip).

Memory bandwidth scaling follows different trajectories across accelerator designs. GPU architectures scale bandwidth by adding memory channels, reaching on the order of 1 TB/s in mainstream products and a few TB/s in high-end systems. TPU-class designs achieve their bandwidth efficiency through systolic array dataflow and aggressive on-chip reuse, often trading flexibility for efficiency on dense tensor kernels. Mobile SoC designs face the tightest constraints, delivering on the order of hundreds of GB/s of unified memory bandwidth within a few-watt power envelope, which demands careful workload scheduling and thermal management.

HBM provides far higher bandwidth than commodity DDR memory, but at substantially higher cost and packaging complexity. High-bandwidth accelerators therefore trade higher memory-system cost for higher sustained performance on bandwidth-bound workloads. Edge accelerators often sacrifice bandwidth to meet tight cost and power targets while maintaining sufficient performance for inference workloads.

These bandwidth characteristics directly influence deployment decisions: cloud training prioritizes raw bandwidth for maximum model capacity, edge inference optimizes bandwidth efficiency for energy constraints, and mobile deployment balances bandwidth with cost limitations. Beyond the accelerator’s internal memory system, host-device traffic can introduce another bottleneck. When an input must cross a 64 GB/s PCIe link before using an accelerator with 2 TB/s of HBM bandwidth, the interface has roughly $30\times$ less bandwidth and can dominate small, frequent transfers (NVIDIA

Corporation 2020a, 2020c). Resident data and direct storage or network paths need not traverse that host link first.

11.5.5 Host-accelerator communication

Machine learning accelerators, such as GPUs and TPUs, achieve high computational throughput through parallel execution. However, their efficiency is often constrained by host-accelerator data movement between the CPU and accelerator memory. Compared to many traditional workloads that keep most data within a single memory domain, AI workloads can require frequent transfers between CPU memory and accelerator memory, introducing latency, consuming bandwidth, and affecting overall performance.

Figure 11.13 traces the host-accelerator sequence for a GPU as the concrete accelerator (its “Memory for GPU” lane is the accelerator’s memory). Before computation begins, data is copied from CPU memory to the accelerator’s memory (step 1). The CPU then issues execution instructions (step 2), and the accelerator processes the data in parallel (step 3). Once computation completes, the accelerator writes its output to accelerator memory (the “Store results” arrow), and that result is copied back to the CPU (step 4). Consider the latency cost at every arrow: each transfer represents a potential bottleneck that must be managed to optimize end-to-end performance.




Figure 11.13: Host-Accelerator Data Transfer Sequence: Four-step execution flow between CPU host and GPU device: (1) copying input data over PCIe, (2) issuing kernel instructions, (3) GPU compute, and (4) transferring results back to host memory. Low PCIe bandwidth and host dispatch latencies can easily dominate runtime unless kernel compute time exceeds data transfer overhead.


The key challenges in host-accelerator data movement include latency, bandwidth constraints, and synchronization overheads. The efficiency of ML accelerators depends as much on a continuous data supply as on raw computational power. Even high-performance GPUs and TPUs remain underutilized if data transfers are inefficient. Host and accelerator memory exist as separate domains, requiring explicit transfers over interconnects such as PCIe, NVLink, or proprietary links. Ineffective data movement causes execution stalls, making transfer optimization a priority.

11.5.5.1 Node-level interconnect topology

Optimizing data movement requires understanding the physical topology of the compute node. A typical AI server is not a flat mesh of connected devices but a hierarchy of bandwidths that tapers as data moves further from the chip.

At node level, three links define the bandwidth taper:

1. **Device-device interconnect (NVLink/Infinity Fabric):** Modern multi-GPU nodes use specialized high-speed bridges like NVLink²⁹ to connect accelerators directly, bypassing the host CPU. Bandwidth ranges from 600 GB/s to 900 GB/s per GPU (NVIDIA Corporation 2020c; Choquette 2023). This link matters whenever tensors must move between accelerators within one server, including model partitioning, activation exchange, and gradient synchronization



Bandwidth tapers steeply as data moves farther from the accelerator.

²⁹ | **NVLink (NVIDIA Link):** This direct GPU-to-GPU interconnect exists to keep accelerator-to-accelerator traffic off PCIe when tensors must move inside a server. Its 600–900 GB/s aggregate bidirectional bandwidth is roughly 300–450 GB/s per direction under balanced full-duplex traffic, several times the directional bandwidth of a standard PCIe link, so workloads with frequent cross-device tensor movement can remain in the fast part of the bandwidth taper (NVIDIA Corporation 2020c; Choquette 2023). The training algorithms that determine how much tensor traffic must cross this link are developed later; the hardware fact needed here is the bandwidth gap.

during training. The hardware lesson for this chapter is the boundary: traffic that stays on the accelerator fabric is far cheaper than traffic that falls back through the host.

2. **Host-device interconnect (PCIe):** The link between the CPU and the accelerator. Bandwidth ranges from 32 to 64 GB/s (PCIe Gen4/Gen5). Host-staged training data can bottleneck here, although resident data, direct storage or network I/O, and coherent accelerator links can use other paths. Multi-GPU servers may also provide multiple PCIe links rather than a single shared 64 GB/s path.
3. **Network interface card:** The link from the host to the outside world, connecting to other nodes. Bandwidth ranges from 25 to 50 GB/s (200 Gb/s to 400 Gb/s Ethernet/InfiniBand³⁰). This interconnect is the first step from single-node hardware reasoning into the next frontier: scale. Here, the point is that leaving the node moves traffic onto a much narrower and higher-latency path.

These three levels produce a characteristic bandwidth taper:

$$\begin{aligned} \text{HBM (3350 GB/s)} &\gg \text{NVLink (900 GB/s)} \\ &\gg \text{PCIe (64 GB/s)} \gg \text{Network (50 GB/s)} \end{aligned}$$

System efficiency depends on keeping data on the fastest suitable path. PCIe and network links provide substantially less bandwidth and higher latency than on-package memory, so placement and scheduling should prevent avoidable host and network crossings.

Offloading computation to specialized hardware introduces bus transfer overheads between host memory and accelerator VRAM.³¹ Trace the numbered steps in figure 11.13, following the sequence from initial PCIe memory copy to command launch, kernel execution, and result retrieval.

Latency and bandwidth limitations directly impact AI workloads. PCIe-class host interconnects are typically much slower than an accelerator's on-package high-bandwidth memory, so large transfers can become bottlenecks, particularly in deep learning tasks. Synchronization overheads compound this problem when computation must wait for data transfers to complete. Efficient scheduling and overlapping transfers with execution are necessary to mitigate these inefficiencies.

11.5.5.2 Transfer optimization

The bandwidth taper described in this section creates a clear optimization hierarchy. Practitioners have two complementary strategies for mitigating transfer overheads: asynchronous data movement and unified memory abstraction.

DMA engines execute bulk transfers after software configures them, avoiding CPU-managed copying of every byte. When separate engines and buffers are available, the next batch can move while the accelerator computes on the current one. This overlap helps only when dependencies permit concurrency and neither path contends for the same limiting resource.

Unified Memory provides the second strategy, offering a single address space accessible by both CPU and accelerator. Rather than requiring explicit copies, the runtime migrates memory pages on demand when either processor accesses them. The programming model simplifies dramatically (a single `malloc` replaces complex staging logic), but introduces performance unpredictability. Page migrations triggered by access patterns can cause latency spikes, and small or scattered accesses may thrash pages back and forth across the interconnect. For this reason, production training workloads typically use explicit DMA-based transfers for predictable performance, while Unified Memory finds its niche in prototyping and workloads where development speed outweighs absolute throughput.

These overheads shape how model operations map to hardware. Convolution commonly exposes spatial reuse that maps well to tiling. Dense attention also has regular matrix structure, but its intermediate state and sequence-length scaling create different capacity and bandwidth pressure. Embeddings, sparse experts, and KV-cache reads add still other access patterns.

11.5.6 Model memory pressure

Model architecture determines which memory term binds. While multilayer perceptrons (MLPs), convolutional neural networks (CNNs), and transformer networks each require large parameter sets, their distinct access patterns create different pressure on weights, activations, bandwidth, and host transfers, so each demands a different accelerator optimization strategy.

³⁰ | **InfiniBand:** Its key feature for multi-node scaling is RDMA (Remote Direct Memory Access). With a GPU-aware path such as GPUDirect RDMA and registered GPU memory, a network adapter can transfer data directly between GPUs across nodes without staging it through host memory. This reduces host involvement and protocol overhead, so scaling is more often constrained by physical bandwidth, topology, and collective implementation rather than CPU-managed packet processing.

³¹ | **DMA (direct memory access):** A dedicated hardware unit that manages the data copy (step 1) without direct CPU management, freeing the CPU to immediately issue computation commands (step 2). This concurrency is critical: without it, the accelerator can idle between compute batches, especially when host-to-device movement is on the critical path.

The lighthouse models introduced in table 1.5 ground this analysis: ResNet-50 represents CNN workloads with high spatial reuse, GPT-2/Llama exemplifies transformer memory pressure, DLRM illustrates sparse embedding lookups that stress memory systems differently than dense operations, and MobileNetV2 demonstrates efficiency-optimized architectures with depthwise convolutions. Their different memory characteristics reveal how workload structure translates to hardware utilization.

11.5.6.1 Multilayer perceptrons

An MLP layer is a dense matrix multiplication followed by optional bias, activation, and normalization. Its weight traffic is expensive at small batch sizes, while larger batches reuse each weight across more examples and can make the matrix multiplication compute bound.

From a memory perspective, MLPs rely on large, dense weight matrices that frequently exceed on-chip SRAM capacity, necessitating accesses to device-local DRAM. PCIe or another host-device interconnect becomes part of the critical path only when inputs are staged from the host or weights and activations are offloaded beyond device memory.

MLPs have regular access patterns that support tiled prefetching and streaming. Accelerators stage weight and activation tiles in on-chip SRAM and overlap movement with arithmetic when dependencies permit. Delivered performance then depends on batch size, tensor shape, precision, and whether the working set remains in device memory (Chen et al. 2017).

11.5.6.2 Convolutional neural networks

CNNs are widely used in image processing and computer vision tasks. Unlike MLPs, which require dense matrix multiplications, CNNs process input feature maps using small filter kernels that slide across the image. This localized computation structure results in high spatial data reuse, where the same input pixels contribute to multiple convolutions.

CNN accelerators benefit from on-chip memory optimizations, as convolution filters exhibit extensive reuse, allowing weights to be stored in fast local SRAM instead of repeatedly accessing device DRAM. Activation maps require careful management due to their size, so CNN accelerators divide them into tiles that fit within on-chip buffers. This reduces device-memory traffic and improves efficiency (Chen et al. 2017).

CNNs can expose substantial weight and activation reuse, but intermediate feature maps may still dominate traffic. Tiling and on-chip buffering reduce device-memory movement; host-device transfer is a separate concern that matters when inputs or outputs are not resident. Eyeriss introduced a row-stationary dataflow to exploit convolutional reuse (Y.-H. Chen et al. 2016). In the broader taxonomy of section 11.8, it is a hybrid that keeps selected rows and partial sums local.

11.5.6.3 Transformer networks

The transformer architectures introduced in section 6.6 have become the dominant architecture for natural language processing and are increasingly used in other domains such as vision and speech recognition. Unlike CNNs, which rely on local computations, transformers perform global attention³² mechanisms, where each token in an input sequence can interact with all other tokens.

These models are particularly challenging for accelerators because global token interaction creates large attention state while GPT-3-scale language models (Brown et al. 2020) can exceed on-chip memory capacity through sheer parameter count. As a result, frequent movement between HBM, caches, and compute units creates substantial latency and bandwidth pressure. If the model spills beyond accelerator memory or uses host offload, PCIe or NVLink transfers add another bottleneck. Unified Memory architectures can mitigate some programming complexity by handling movement between host and device memory at runtime, but they introduce additional latency when page migrations occur unpredictably. These pressures make high-bandwidth memory, tensor tiling, and memory partitioning central accelerator design concerns for transformer workloads.

Attention tiling, KV caching, and tensor layouts reduce or reorganize memory traffic. When transformer training partitions tensors across GPUs within a node, NVLink can provide more bandwidth than PCIe for the required exchanges; it offers no benefit if the workload does not cross that link. Asynchronous transfers can overlap movement with computation when the schedule and resources permit (Narayanan et al. 2021).

³² | **Attention mechanism:** Bahdanau, Cho, and Bengio introduced attention for sequence-to-sequence models; Transformer self-attention later allowed each token to attend to every token in the sequence (Bahdanau et al. 2015; Vaswani et al. 2017); materializing full attention scores for a sequence of length S requires an $S \times S$ matrix, producing quadratic intermediate storage. FlashAttention tiles the computation to avoid materializing that full matrix in HBM (T. Dao et al. 2022). The inference KV cache is separate: its capacity grows linearly with sequence length per layer, although each decode step reads an increasing cache (see section 13.8.3).

11.5.7 Accelerator design implications

The diverse memory requirements of MLPs, CNNs, and transformers highlight the need for workload-specific accelerator design. Table 11.11 reveals how memory access patterns vary dramatically across model types.

Table 11.11: ML Model Memory Access: Different machine learning models exhibit distinct memory access patterns and bottlenecks due to variations in weight size, activation reuse, and data movement. Standard dense transformers demand high bandwidth and capacity because large weights, KV caches, and attention traffic dominate memory pressure; sparse MoE or pruned variants add further routing and sparsity considerations. CNNs benefit from spatial locality and high activation reuse, reducing memory pressure.

Model type	Weight size	Activation reuse	Memory access pattern	Primary Bottleneck
MLP (Dense)	Large, dense	Low	Regular, sequential (streamed)	Bandwidth (off-chip)
CNN	Small, reused	High	Spatial locality	Feature map movement
Transformer	Large, usually dense; sparse in MoE/pruned variants	Low-to-medium	Mostly regular GEMM plus KV-cache/attention traffic	Memory capacity + bandwidth

Each model type creates different pressure. Small-batch MLP inference may stream large resident weight matrices from device memory, whereas larger batches increase reuse. CNNs can exploit spatial tiling and filter reuse. Transformer memory pressure depends on phase: training retains activations and exchanges partitioned tensors, while autoregressive inference repeatedly reads weights and an expanding KV cache.

To address these challenges, modern AI accelerators incorporate multi-tier memory hierarchies that balance speed, capacity, and energy efficiency. On-chip SRAM caches and scratchpad memories store frequently accessed data, while high-bandwidth external memory provides scalability for large models. Efficient interconnects, such as NVLink, help alleviate host-accelerator transfer bottlenecks, particularly in transformer workloads where memory movement constraints can dominate execution time.

Memory efficiency can matter as much as raw compute power. Representative energy studies place DRAM access far above simple on-chip arithmetic, while model architectures differ in reuse, capacity, and communication demand. The remaining question is quantitative: on a given accelerator, is a specific kernel limited by arithmetic or by the relevant memory path? The roofline model answers that question before engineers optimize the wrong resource.


11.6 Roofline Model

The Roofline Model answers this question by plotting arithmetic intensity against attainable performance, revealing whether an operation meets a compute ceiling or a memory-bandwidth ceiling. Peak FLOP/s remains one input, but the model combines it with workload traffic and platform bandwidth to produce a more useful bound.

The Roofline Model³³ (Williams et al. 2009) provides the standard framework for understanding whether workloads are compute bound or memory bound, directly connecting the memory wall discussion to practical performance analysis. This model enables quantitative reasoning about accelerator utilization and guides optimization decisions.

Performance is bounded by two ceilings, as equation 11.2 formalizes. Here, attainable performance R_{attain} and peak compute R_{peak} are in FLOP/s (often reported as TFLOP/s), peak bandwidth BW is in bytes/s (often TB/s), and arithmetic intensity I is in FLOP/byte:

$$R_{\text{attain}} = \min(R_{\text{peak}}, \text{BW} \times I) \quad (11.2)$$



Low arithmetic intensity pins the workload in the memory-bound regime.

³³ | **Roofline model:** Introduced by Williams et al. (2009) at UC Berkeley, building on earlier I/O complexity work from the 1980s. Their specific contribution was making the compute vs. bandwidth trade-off *visual* and *actionable*: the characteristic roofline plot immediately reveals whether a kernel is compute bound (hitting the flat ceiling) or memory bound (hitting the sloped bandwidth line) and quantifies the gap to hardware limits. A kernel operating at only 50 percent of its ceiling has a clear $2\times$ utilization gap to close, making this the standard diagnostic tool for accelerator optimization.

Definition 11.5: Arithmetic intensity

Arithmetic Intensity is the ratio of operations to bytes transferred across the memory interface being modeled (FLOP/byte). Together with that level's bandwidth and the relevant compute peak, it determines which idealized roofline ceiling is lower.

1. **Significance:** The threshold between the two ceilings is the ridge point $R_{\text{peak}}/\text{BW}$. For an A100 (312 TFLOP/s FP16/BF16, 2.04 TB/s), it is roughly 153 FLOP/byte. Under the stated traffic models, a large well-tiled 1024×1024 multiply reaches about 341.3 FLOP/byte, while standalone ReLU is about 0.125. They therefore meet different ceilings on the same hardware.
2. **Distinction:** Unlike total FLOPs (a count of operations), arithmetic intensity is a ratio that characterizes the shape of a workload's hardware demand. Two kernels with identical FLOPs but different memory access patterns have different arithmetic intensities and will be bottlenecked by different hardware resources.
3. **Common pitfall:** A frequent misconception is that arithmetic intensity is a fixed property of an operation. In practice, it depends on implementation details: a naive matrix-multiply that reloads operands from DRAM for each output element has low arithmetic intensity; a blocked (tiled) implementation that reuses data from fast SRAM achieves high arithmetic intensity—the same mathematical operation, orders of magnitude apart in hardware efficiency.

With O as the dimensionless operation count and D_{vol} as data volume in bytes, equation 11.3 makes the definition operational.

$$I = \frac{O}{D_{\text{vol}}} \quad (11.3)$$

The roofline visualization shows performance (TFLOP/s) on the vertical axis and arithmetic intensity (FLOP/byte) on the horizontal axis. At low arithmetic intensity, performance increases linearly with intensity (memory-bound region). Above the ridge point, performance saturates at peak compute (compute-bound region). Section A.5.1 maps each regime to the optimizations that pay off and the ones that waste effort, so that classifying a workload as memory-bound or compute-bound tells the engineer directly whether a faster accelerator or more bandwidth is the binding investment.

11.6.1 Hardware ridge points

The ridge point, the hardware balance I_{ridge} established in section 11.5.2, is the arithmetic-intensity threshold at which an accelerator turns from memory-bound to compute-bound. Table 11.12 quantifies how different accelerators exhibit distinct characteristics based on their compute-to-bandwidth ratios:

Table 11.12: Hardware Ridge Points: Vendor-published FP16 peak throughput and memory bandwidth for three NVIDIA generations, and the ridge point each pair implies. The ridge rises from V100 to H100 because compute throughput grows faster than bandwidth; these values are hardware-specific rather than universal generation averages.

NVIDIA Accelerator	Peak FP16	Bandwidth	Ridge Point
V100 (2017)	125 TFLOP/s	0.9 TB/s	138.9 FLOP/byte
A100 (2020)	312 TFLOP/s	2.04 TB/s	153 FLOP/byte
H100 (2022)	989 TFLOP/s	3.35 TB/s	295.2 FLOP/byte

These ridge points show that later tensor paths require more reuse to reach their arithmetic ceiling. A ridge-point comparison makes that trend concrete.

Napkin Math 11.3: The utilization gap

Problem: Why is it harder to get 100 percent utilization on an H100 than a V100?

Metric: The ridge point $I_{\text{ridge}} = R_{\text{peak}}/\text{BW}$ (FLOP/byte) from section 11.5.2: how many math operations the hardware must perform for every byte of data loaded to keep the compute units busy.

Evolution:

- **V100 (2017):** 125 TFLOP/s / 0.9 TB/s $\approx$ 138.9 FLOP/byte.
- **A100 (2020):** 312 TFLOP/s / 2.04 TB/s $\approx$ 153 FLOP/byte.
- **H100 (2022):** 989 TFLOP/s / 3.35 TB/s $\approx$ 295.2 FLOP/byte.

Result: The intensity needed to reach the tensor peak has increased. Under this simple HBM roofline, a kernel with $I = 200$ FLOP/byte lies above the A100 ridge but below the H100 ridge. If its intensity and implementation remain unchanged, its idealized gain is then closer to the $1.6\times$ bandwidth ratio than the $3.2\times$ peak-compute ratio.

Systems insight: A standalone ReLU under a one-read, one-write FP32 traffic model performs about 1 operation per 8 bytes, far below the H100 tensor ridge. A well-tiled 1024×1024 dense multiply reaches about 341.3 FLOP/byte under its idealized traffic model. Fusion is valuable for low-intensity chains because it can remove intermediate round-trips; it is not automatically the best optimization for every kernel.

Depthwise convolution, embedding lookup, LayerNorm, and softmax are useful low-intensity reference points because they spend more time moving bytes than doing arithmetic. Table 11.13 maps common neural network operations to the Roofline Model.

Table 11.13: Operations on the Roofline: Neural network layers span a wide range of arithmetic intensities. Large, well-tiled convolutions and batched GEMMs can be compute-bound, while MobileNet depthwise layers, attention softmax, normalization, and DLRM embeddings are often memory-bound. To see how these intensity values translate into real performance predictions, a transformer layer provides a complete arithmetic-intensity calculation across its major sub-operations.

Operation	Arithmetic Intensity	Classification	Lighthouse example
Conv2D (Dense)	50–200 FLOP/byte	Straddles ridge; high-reuse cases compute-bound	ResNet-50
Dense MatMul (large batch, well-tiled)	64–256+ FLOP/byte	Often compute-bound at large batch	GPT-2 (batched projections)
Depthwise conv	10–20 FLOP/byte	Memory-bound	MobileNet
Attention Softmax	Convention-dependent	Memory-bound	GPT-2 (Generation)
LayerNorm	1–2 FLOP/byte	Memory-bound	GPT-2/Llama
Embedding lookup	< 1 FLOP/byte	Memory-bound	DLRM

Napkin Math 11.4: Transformer layer analysis

For a transformer with hidden_dim = 768, batch = 32, seq = 512:

Attention QKV projection:

- FLOPs: $2 \times 3 \times 32 \times 512 \times 768 \times 768 = 58$ GFLOP
- Bytes: (input + weights + output) = $(32 \times 512 \times 768 + 3 \times 768 \times 768 + 32 \times 512 \times 768 \times 3) \times 2 \approx 104.2$ MB
- AI = 58 GFLOP / 104.2 MB = 556.4 FLOP/byte, which is compute bound on A100 (above 153 FLOP/byte threshold)

Softmax:

- FLOPs: $32 \times 12 \times 512 \times 512 \times 3 \approx 302$ MFLOP (exp, sum, div)

- Bytes: $32 \times 12 \times 512 \times 512 \times 2 \times 2 = 402.7$ MB
- AI = $302 \text{ MFLOP} / 402.7 \text{ MB} = 0.75 \text{ FLOP/byte}$ under this simple operation count, which is memory-bound. Counts that weight exponentials and division as multiple operations produce a larger numerical intensity without changing the bottleneck classification.

Systems insight: This analysis explains why FlashAttention focuses on reducing memory traffic in attention rather than reducing FLOPs.

These classifications directly inform optimization strategy. Memory-bound operations benefit from reducing data movement through operator fusion, using reduced precision (FP16, INT8), and increasing arithmetic intensity through algorithmic changes like FlashAttention. Compute-bound operations, by contrast, benefit from maximizing hardware utilization through batching and parallelism, exploiting Tensor Cores and specialized compute units, and optimizing compute efficiency through tiling and scheduling.

11.6.2 Calculating memory bandwidth bounds

The Roofline Model's memory-bound region is determined by the peak memory bandwidth. For an operation to achieve throughput R_{ops} (FLOP/s, often expressed in TFLOP/s) in the memory-bound regime, equation 11.4 gives the required bandwidth:

$$\text{BW}_{\text{req}} = \frac{R_{\text{ops}}}{I} \text{ bytes/s} \quad (11.4)$$

When required bandwidth exceeds peak bandwidth, performance is capped according to equation 11.5. Here R_{ops} and R_{attain} are in FLOP/s and I is in FLOP/byte.

$$R_{\text{attain}} = \text{BW} \times I \quad (11.5)$$

A convolution layer provides the compute-bound contrast.



Napkin Math 11.5: Convolutional layer analysis

Consider a Conv2D layer with input shape (batch = 32, channels = 128, height = 56, width = 56), output channels = 256, kernel size 3×3 on an A100 GPU, which provides the compute-bound contrast:

Computational requirements:

- Output size: $32 \times 256 \times 56 \times 56 = 25.7\text{M}$ elements
- FLOPs per output: $128 \times 3 \times 3 \times 2 = 2,304$ (multiply-add)
- Total FLOPs: $25.7\text{M} \times 2,304 = 59.2 \text{ GFLOP}$

Memory traffic analysis:

- Input: $32 \times 128 \times 56 \times 56 \times 2 = 25.7 \text{ MB}$ (FP16)
- Weights: $256 \times 128 \times 3 \times 3 \times 2 \approx 0.6 \text{ MB}$ (FP16)
- Output: $32 \times 256 \times 56 \times 56 \times 2 = 51.4 \text{ MB}$ (FP16)
- Total: 77.7 MB

Arithmetic intensity: $I = 59.2 \text{ GFLOP} / 77.7 \text{ MB} = 762.2 \text{ FLOP/byte}$

Systems insight: This is well above A100's ridge point of 153 FLOP/byte, so the simple Roofline Model places the operation in the compute-limited region. Peak throughput of $\sim 312 \text{ TFLOP/s}$ (FP16 with Tensor Cores) is therefore the ceiling, but realized performance still depends on tiling, occupancy, instruction mix, and library efficiency.

The convolutional layer's high arithmetic intensity arises from its weight reuse pattern: the same 3×3 kernel is applied across all spatial locations, amortizing the cost of loading weights across

millions of output computations. This is the architectural pattern that makes CNNs so efficient on modern accelerators.

Napkin Math 11.6: Dense layer analysis

Consider a fully connected layer: input (batch = 32, features = 2048) → output (batch = 32, features = 2048) on the same A100:

Computational requirements:

- Matrix multiply: $(32 \times 2048) \times (2048 \times 2048)$
- Total FLOPs: $2 \times 32 \times 2048 \times 2048 = 268.4$ MFLOP

Memory traffic analysis:

- Input: $32 \times 2048 \times 2 = 131.1$ KB (FP16)
- Weights: $2048 \times 2048 \times 2 = 8.4$ MB (FP16)
- Output: $32 \times 2048 \times 2 = 131.1$ KB (FP16)
- Total: 8.7 MB

Arithmetic intensity: $I = 268.4 \text{ MFLOP} / 8.7 \text{ MB} = 31 \text{ FLOP/byte}$

This intensity sits below A100's ridge point of 153 FLOP/byte, making this operation memory-bound. Attainable performance: $R_{\text{attain}} = 2,039 \text{ GB/s} \times 31 \text{ FLOP/byte} = 63.3 \text{ TFLOP/s}$

Systems insight: Delivered throughput reaches only 20.3 percent of peak compute capability, demonstrating the memory wall effect for small batch sizes.

However, not all layers in a neural network exhibit this favorable profile. The fully connected (dense) layers that typically appear at the end of classification networks, or as the projection layers in transformers, have different arithmetic intensity characteristics. A dense layer provides the memory-bound contrast needed to predict where bottlenecks will occur in end-to-end model execution.

The dense layer's lower arithmetic intensity stems from limited weight reuse: each weight is reused across the batch but lacks the additional spatial reuse of convolutional filters, so small-batch dense layers have much lower arithmetic intensity than convolutions. This difference explains why transformer inference (dominated by dense projections) is typically memory bound while CNN inference can be compute bound.

Napkin Math 11.7: LayerNorm analysis

LayerNorm with input shape (batch = 32, seq = 512, hidden = 768):

Computational requirements:

- Elements: $32 \times 512 \times 768 = 12.6\text{M}$
- Approximate operations per element: mean reduction (1 ADD), variance (2 ADD, 1 MUL), normalization and affine transform (2 ADD, 2 MUL) ≈ 8 ; reciprocal-square-root overhead is amortized across each hidden vector
- Total FLOPs: $12.6\text{M} \times 8 = 100.7$ MFLOP

Memory traffic:

- Input: $12.6\text{M} \times 2 = 25.2$ MB
- Parameters (scale, bias): $768 \times 2 \times 2 = 3.1$ KB (negligible)
- Output: $12.6\text{M} \times 2 = 25.2$ MB
- Total: 50.3 MB

Arithmetic intensity: $I = 100.7 \text{ MFLOP} / 50.3 \text{ MB} = 2.0 \text{ FLOP/byte}$

This arithmetic intensity sits severely below the A100 ridge point (77× below). Performance is limited to: $R_{\text{attain}} = 2039 \text{ GB/s} \times 2.0 \text{ FLOP/byte} = 4.1 \text{ TFLOP/s}$

Systems insight: The simple HBM roofline bounds this standalone LayerNorm at about 1 percent of A100’s tensor peak. Actual latency also depends on fusion, reduction strategy, launch overhead, and which memory level supplies the data.

Standalone normalization performs little arithmetic relative to the bytes it touches, as the LayerNorm example shows. Within that isolated pass, each element is read and written with little reuse; fusion with adjacent operators can remove some of those round-trips.

11.6.3 Optimization by intensity regime

The roofline analysis directly informs optimization priorities, summarized in table 11.14.

Table 11.14: Optimization by Arithmetic Intensity Regime: Roofline position determines whether an optimization should chase compute utilization, memory-traffic reduction, or complete elimination of memory round-trips. The lower the arithmetic intensity, the more valuable fusion and data-movement avoidance become.

Intensity regime	Typical operations	Optimization priority	Common techniques	Expected impact
High AI (> 200 FLOP/byte)	Large convolutions	Maximize compute utilization.	Tensor Cores, thread-block tuning, and high occupancy.	Raises sustained compute utilization when the kernel already clears the ridge.
Medium AI (20–200 FLOP/byte)	Medium-sized dense layers	Balance compute and memory optimization.	Larger batches, register tiling, and fusion with adjacent operations.	Can move the binding bottleneck between memory and compute.
Low AI (< 20 FLOP/byte)	Small dense layers and element-wise operations	Reduce memory traffic.	Aggressive operator fusion, reduced precision (FP16 $\rightarrow$ INT8), and algorithmic changes.	Reduces external-memory traffic when fusion or layout changes are legal.
Very low AI (< 2 FLOP/byte)	Normalization layers and activation functions	Eliminate memory round-trips.	Fuse with adjacent operations and use in-place computation where legal.	Removes intermediate round-trips when adjacent operations can be fused.

For low-AI operations, operator fusion is often the decisive optimization: LayerNorm combined with the Gaussian Error Linear Unit (GELU), for example, can become a single fused kernel. One of the most accessible levers for moving an operation up and right on the roofline is batching.



Napkin Math 11.8: Batch size and arithmetic intensity

Increasing batch size improves arithmetic intensity for matrix operations by amortizing weight loading. Equation 11.6 formalizes this relationship for an idealized FP16/BF16 dense layer $(B \times M) \times (M \times N)$, counting one read each of the input and weights and one write of the output, where B is batch size and M and N are the input and output dimensions:

$$I = \frac{2BMN}{2BM + 2MN + 2BN} \approx \frac{2BMN}{2MN} = B \quad (\text{when } 2MN \gg 2B(M + N)) \quad (11.6)$$

Example: Dense layer with $M=N=2048$ (FP16)

- Batch = 1: AI ≈ 1 FLOP/byte (memory bound)
- Batch = 32: AI ≈ 31 FLOP/byte (memory bound)
- Batch = 256: AI ≈ 204.8 FLOP/byte (compute bound on A100)

Systems insight: This explains why batching can produce large throughput improvements in production inference systems, as MLPerf Inference, the standardized benchmark suite covered in chapter 12, demonstrates by separating bulk-throughput runs from latency-constrained serving runs (Reddi et al. 2019).

The batch size analysis reveals why inference serving systems are designed around batching: it changes the arithmetic intensity regime of memory-bound workloads. However, batching introduces

latency trade-offs, since requests must wait in a queue until a batch forms. This tension between throughput (favoring large batches) and latency (favoring small batches) is a central challenge in ML serving systems, explored in depth in section 13.7.3.

When service latency limits batching, batch-1 or small-batch LLM decode often retains low weight reuse and low arithmetic intensity. The next worked model quantifies that particular regime rather than treating all LLM inference as identical.

Napkin Math 11.9: The throughput ceiling

Problem: A batch-1 GPT-2 calculation makes this bandwidth ceiling concrete: what is the maximum possible utilization of an NVIDIA A100 when running GPT-2 inference (batch size 1)?

The hardware constraints (the denominators)

- **Peak compute:** 312 TFLOP/s (FP16 Tensor Core).
- **Peak bandwidth:** 2.04 TB/s (HBM2e).
- **Ridge point** ($R_{\text{peak}}/\text{BW}$): $312 \text{ TFLOP/s} / 2.04 \text{ TB/s} = 153 \text{ FLOP/byte}$ (for FP16 Tensor Core).

Interpretation: Saturating this chip at FP16 precision requires 153 FLOP/byte operations for every byte loaded. The ridge point varies by precision: FP32 operations (19.5 TFLOP/s peak) have a ridge point of only $\sim 9.6 \text{ FLOP/byte}$.

The workload characteristics (the numerator)

- **Model:** GPT-2 XL (1.5 billion parameters).
- **Operation:** Autoregressive generation (1 token at a time).
- **Data movement:** Must load all weights (3 GB @ FP16) for every token.
- **Compute:** Vector-Matrix multiplication. $2 \times \text{Params} \approx 3 \text{ GFLOP}$.
- **Arithmetic intensity:** $3 \text{ GFLOP} / 3 \text{ GB} = 1 \text{ FLOP/byte}$

The prediction (iron law)

Since Actual Intensity (1) $\ll$ Ridge Point (153 FLOP/byte), the system is bandwidth bound.

- **Maximum throughput:** $1 \text{ FLOP/byte} \times 2.04 \text{ TB/s} = 2.04 \text{ TFLOP/s}$.
- **Utilization ceiling:** $2.04 \text{ TFLOP/s (Actual)} / 312 \text{ TFLOP/s (Peak)} \approx 0.7\%$

Systems insight: Without batching, a \$15,000 GPU runs at less than 1 percent compute utilization in this weight-streaming decode model. Batching and weight quantization address that gap; key-value caching instead avoids recomputing prior attention states.

Through this derivation, the Roofline Model provides a diagnostic framework for identifying whether operations are compute bound or memory bound. Knowing that a workload is memory bound at 0.7 percent utilization is only the first step; the next challenge is translating this diagnosis into efficient execution plans that exploit accelerator architectures.

11.7 Hardware Mapping

Consider a 3×3 convolution running on an accelerator tile. The mathematical operation is fixed, but the execution plan is not. One schedule can keep a filter in local registers while many output pixels stream past it; another can advance pixel by pixel and reload the same filter values repeatedly from a slower memory tier. Both schedules compute the same tensor. Only one turns the reuse in the convolution into real bandwidth savings.

Definition 11.6: Mapping in AI acceleration

Mapping in AI Acceleration is the accelerator-compiler process of binding the Logical Computation Graph to the Physical Hardware Topology by deciding which operations execute on which processing elements, which data resides in which memory tier, and in what temporal order.

1. **Significance:** Within the D·A·M taxonomy, mapping is a machine-axis decision that helps determine how closely an operation approaches the roofline bound $\min(R_{\text{peak}}, \text{BW} \times I)$. A poor tiling choice that forces unnecessary DRAM accesses can reduce effective arithmetic intensity enough to move a nominally compute-bound operation into the bandwidth-bound regime.
2. **Distinction:** Traditional compilation often emphasizes instruction selection, register allocation, and scheduling for general-purpose processors. Accelerator mapping gives equal prominence to spatial placement and explicit data movement across a Dataflow Architecture. The reason is physical: in the 45 nm comparison introduced earlier, off-chip DRAM access consumed roughly $200\times$ the energy of a local integer multiply-accumulate, although the ratio depends on the technology and operation.
3. **Common pitfall:** A frequent misconception is that mapping is automatically handled by frameworks. For general GPU workloads, compilers like Accelerated Linear Algebra (XLA) can find strong mappings for common kernels; for specialized accelerators (systolic arrays, custom ASICs), compiler-generated mappings may still lag hand-tuned schedules because the compiler's search space is limited by the time budget at compilation.

The convolution example exposes the three decisions that recur throughout accelerator compilation. Placement assigns the multiply-accumulate work to processing elements so parallelism does not turn into idle time or interconnect congestion. Allocation keeps weights, activations, and partial sums in the memory tier where their next use will occur, rather than letting reuse spill back to DRAM. Scheduling orders loops and kernels so the chosen placement and allocation remain valid over time. A poor choice in any one dimension can collapse a high-arithmetic-intensity operation back into a bandwidth-bound execution. In practice, these choices are too coupled for developers to manage by hand at model scale, which is why systems such as XLA, TVM, and TensorRT lower high-level models and search or select execution plans within their compile-time and hardware budgets. Section 11.9 examines that compiler support in detail.

11.7.1 Placement and allocation

Translating a model's computational graph into efficient hardware execution requires solving two tightly coupled problems. Computation placement determines which operations run on which processing elements, balancing parallelism against communication costs. Memory allocation determines where data resides within the memory hierarchy, trading capacity against access latency. These two decisions interact: placing operations on distant processing elements increases the memory bandwidth required to shuttle data between them, while allocating data to fast but small on-chip memory limits which operations can execute concurrently. Getting either wrong leaves thousands of processing elements idle or starved for data.

11.7.1.1 Computation placement

Computation placement is the process of assigning operations to an accelerator's processing elements (PEs) to expose parallelism, limit idle time, and reduce unnecessary data movement. Modern accelerators contain many such resources: the NVIDIA H100 has more than 16,000 CUDA cores and more than 500 Tensor Cores (Choquette 2023), TPUs organize thousands of multiply-accumulate units into systolic arrays (Jouppi et al. 2017), and wafer-scale processors such as Cerebras' CS-2 integrate more than 850,000 cores (Systems 2021). At these scales, placement inefficiencies become measurable because idle cores and redundant transfers waste both time and energy.

The difficulty of placement depends on workload regularity. CNNs and dense transformer operations are structured and tile well, although transformers combine kernels with different shapes

and reuse patterns. Graph Neural Networks (GNNs) are harder to partition because sparse, input-dependent neighborhoods can create load imbalance and scattered communication. Table 11.15 lists the core challenges placement must address across these workload types. Static schedules work well for many fixed-shape dense kernels; runtime-aware placement is useful when shapes, sparsity, routing, or resource availability vary.

Table 11.15: Computation Placement Challenges: Effective neural network deployment requires strategic allocation of computations to processing elements, balancing workload distribution, data movement costs, and hardware constraints to maximize execution efficiency. These challenges guide the design of mapping strategies that optimize resource utilization and minimize communication overhead.

Challenge	Impact on execution	Key Considerations for Placement
Workload imbalance	Some processing elements finish early while others remain overloaded, leading to idle compute resources.	Distribute operations evenly to prevent stalls and ensure full utilization of PEs.
Irregular computation patterns	Sparse graphs, routing, and variable shapes can create nonuniform work that complicates static placement.	Use adaptive placement when workload characteristics vary at run time.
Excessive data movement	Frequent memory transfers introduce latency and increase power consumption.	Keep frequently used data close to the compute units and minimize off-chip memory accesses.
Limited interconnect bandwidth	Poorly placed operations can create congestion, slowing data movement between PEs.	Optimize spatial and temporal placement to reduce communication overhead.
Model-specific execution needs	CNNs, transformers, and GNNs require different execution patterns, making a single placement strategy ineffective.	Tailor placement strategies to match the computational structure of each model type.

Good placement can substantially reduce latency, while poor placement leaves processing elements idle or increases communication. Accelerators therefore combine static mappings for regular kernels with runtime-aware scheduling where workload behavior varies. Placement decisions also interact directly with the next concern: where the data those processing elements need resides in the memory hierarchy.

11.7.1.2 Memory allocation

While computation placement determines where operations execute, memory allocation defines where data resides and how it flows through the hierarchy. The goal is to keep reused data near the processing elements without exceeding faster-tier capacity. GPUs expose global memory, shared memory, caches, and registers that kernels coordinate through tiling (NVIDIA Corporation 2020a). TPUs use on-chip buffers to stage activations and weights for systolic-array execution (figure 11.8) (Jouppi et al. 2017), while wafer-scale processors partition memory and computation to control interconnect traffic (Systems 2021). General-purpose processors and accelerators both use memory hierarchies, but many accelerators expose more placement decisions to compilers and kernels. Poor allocation can therefore impose three related penalties: additional latency, greater energy from higher-tier accesses, and lower throughput when processing elements wait for data.

The severity of these penalties varies by workload. CNNs rely on structured, localized access patterns and benefit from well-defined memory layouts that facilitate predictable reuse (Y.-H. Chen et al. 2016). Transformer models require access to large parameter sets and intermediate activations, making them sensitive to capacity and bandwidth. GNNs add irregular sparse structures that complicate allocation and prefetching.

Capacity is only the first allocation test. A tensor can fit in accelerator memory yet still dominate execution if each kernel rereads it from a distant tier; conversely, tiling can keep a larger working set efficient by preserving reuse locally. The allocation diagnostic is therefore traffic, not fit alone: designers must count how often values cross each memory boundary and keep frequently reused values in the fastest feasible tier. Table 11.16 summarizes these challenges. Systems combine static memory plans for known shapes with dynamic buffer management when shapes vary, while device capacity limits which models fit without partitioning or offload.

Table 11.16: Memory Allocation Challenges: Efficient memory management in AI accelerators balances data access speed with hardware constraints, mitigating performance bottlenecks caused by latency, bandwidth limitations, and irregular data patterns. Complex models such as transformers and graph networks impose variable and demanding memory requirements that amplify these challenges.

Challenge	Impact on Execution	Key Considerations for Allocation
High memory latency	Slow data access delays execution and reduces throughput.	Prioritize placing frequently accessed data in faster memory locations.
Limited on-chip storage	Small local memory constrains the amount of data available near compute units.	Allocate storage efficiently to maximize data availability without exceeding hardware limits.
High off-chip bandwidth demand	Frequent access to external memory increases delays and power consumption.	Reduce unnecessary memory transfers by carefully managing when and how data is moved.
Irregular memory access patterns	Some models require accessing data unpredictably, leading to inefficient memory usage.	Organize memory layout to align with access patterns and minimize unnecessary data movement.
Model-specific memory needs	Different models require different allocation strategies to optimize performance.	Tailor allocation decisions based on the structure and execution characteristics of the workload.

11.7.2 Combinatorial complexity

The small convolution example also explains why hardware mapping becomes a combinatorial search problem. Keeping a filter local improves reuse only if the chosen processing elements have enough nearby storage and if the loop order revisits that filter before eviction. Parallelizing across more processing elements improves throughput only until synchronization and interconnect traffic consume the gain. Table 11.17 lists recurring tensions among placement, allocation, and scheduling. Because changing one decision alters the feasible and useful choices for the others, practical mappers search or optimize them jointly rather than choosing each dimension independently.

These interacting factors define a vast combinatorial design space where small variations in mapping decisions lead to large differences in performance and energy efficiency. Unlike traditional workloads with predictable execution patterns, machine learning models introduce diverse computational structures that require mappings adapted to data reuse, parallelization opportunities, and memory constraints. The search space grows combinatorially, making exhaustive search infeasible. Three sources of variation contribute to this complexity:

Table 11.17: Placement-Allocation-Scheduling Trade-Offs: AI accelerator performance depends on mapping computations to hardware, allocating data to memory tiers, and scheduling execution over time. Careful consideration of these interdependent factors is essential for maximizing throughput and minimizing energy consumption.

Dimension	Placement considerations	Allocation and Scheduling Considerations
Computational granularity	Fine-grained placement enables greater parallelism but increases synchronization overhead.	Coarse-grained scheduling reduces synchronization overhead but may limit flexibility.
Spatial vs. Temporal Mapping	Spatial placement enhances parallel execution but can lead to resource contention and memory congestion.	Temporal scheduling balances resource sharing but may reduce overall throughput.
Memory and Data Locality	Placing data closer to compute units minimizes latency but may reduce overall memory availability.	Allocating data across multiple memory levels increases capacity but introduces higher access costs.
Communication and synchronization	Co-locating compute units reduces communication latency but may introduce contention.	Scheduling synchronization mechanisms mitigates stalls but can introduce additional overhead.
Dataflow and Execution Ordering	Static placement simplifies execution but limits adaptability to workload variations.	Dynamic scheduling improves adaptability but adds scheduling complexity.

11.7.2.1 Ordering computation and execution

Machine learning workloads are often structured as nested loops that iterate over various dimensions of computation. For instance, a matrix multiplication kernel may loop over batch size (B), input features (C_{in}), and output features (C_{out}). The order in which these loops execute has a profound effect on data locality, reuse patterns, and computational efficiency.

Ignoring dependences and equivalent orderings, the number of ways to arrange n_{loops} loops has the toy upper bound:

$$N_{\text{order}} = n_{\text{loops}}!$$

which scales rapidly. A typical convolutional layer may involve up to seven loop dimensions, leading to:

$$7! = 5,040 \text{ possible execution orders.}$$

If each memory level could choose an ordering independently, the corresponding toy upper bound would expand as:

$$(n_{\text{loops}}!)^{N_{\text{mem}}}$$

where N_{mem} is the number of memory hierarchy levels. This rapid expansion shows why execution order optimization matters: poor loop ordering can lead to excessive memory traffic, while an optimized order improves cache utilization (Sze et al. 2017).

11.7.2.2 Parallelization across processing elements

Modern AI accelerators use thousands of processing elements to maximize parallelism, but determining which computations should be parallelized requires careful analysis. Excessive parallelization can introduce synchronization overheads and increased bandwidth demands, while insufficient parallelization leads to underutilized hardware.

Before legality and capacity constraints are applied, the number of ordered ways to select loops for parallel execution has the upper bound:

$$\mathcal{P}_{\text{parallel}} = \frac{n_{\text{loops}}!}{(n_{\text{loops}} - k_{\text{parallel}})!}$$

where n_{loops} is the number of loops, and k_{parallel} is the number selected for parallel execution. For a six-loop computation where three loops are selected, this unconstrained count is:

$$\frac{6!}{(6-3)!} = 120.$$

Even for a single layer, the candidate count can reach hundreds before invalid or equivalent strategies are removed. Each remaining strategy affects data synchronization, memory contention, and overall compute efficiency.

11.7.2.3 Memory placement and data movement

The hierarchical memory structure of AI accelerators introduces additional constraints, as data must be efficiently placed across registers, caches, shared memory, and off-chip DRAM. Data placement impacts latency, bandwidth consumption, and energy efficiency. Frequent access to slow memory creates bottlenecks, while optimized placement reduces costly memory transfers.

If each computational dimension had n independent choices at every memory level, the resulting toy upper bound would be:

$$\mathcal{M}_{\text{placement}} = n^{N_{\text{comp}} \times N_{\text{mem}}}$$

where:


- n = number of placement choices per level,
- N_{comp} = number of computational dimensions,
- N_{mem} = number of memory hierarchy levels.

For a model with:

- $N_{\text{comp}} = 5$ computational dimensions,
- $N_{\text{mem}} = 3$ memory levels,
- $n = 4$ possible placement choices per level,

the number of possible memory allocations is:

$$4^{5 \times 3} = 4^{15} = 1,073,741,824.$$



11.7.2.4 Mapping search space

This unconstrained mapping example exceeds a billion combinations, although many are illegal, coupled, or equivalent. Multiplying the toy counts gives an upper-bound illustration of the mapping search space:

$$\mathcal{S}_{\text{mapping}} = \left(n^{N_{\text{comp}}} \times n_{\text{loops}}! \times \frac{n_{\text{loops}}!}{(n_{\text{loops}} - k_{\text{parallel}})!} \right)^{N_{\text{mem}}}$$

where:

- $n^{N_{\text{comp}}}$ represents memory placement choices,
- $n_{\text{loops}}!$ accounts for computation ordering choices,
- $\frac{n_{\text{loops}}!}{(n_{\text{loops}} - k_{\text{parallel}})!}$ captures parallelization possibilities,
- N_{mem} is the number of memory hierarchy levels.

This equation is not an exact count of legal schedules because choices interact and hardware constraints prune the space. It illustrates why candidate spaces can grow rapidly and why exhaustive search becomes impractical. A concrete example makes the impact of these choices tangible.



Example 11.4: Loop ordering in a small convolution

Consider a convolution applying 16 filters of size 3×3 to an 8×8 single-channel input. The computation can be expressed as five nested loops iterating over output rows (H_{out}), output columns (W_{out}), filter count (C_{out}), filter height (F_h), and filter width (F_w). The $5! = 120$ permutations are an unconstrained upper bound: legal orderings preserve the mathematical result, although floating-point reduction order can change rounding, and they generate different memory traffic.

Ordering A (weight-stationary): Place the filter loops (C_{out}, F_h, F_w) outermost and the spatial loops ($H_{\text{out}}, W_{\text{out}}$) innermost. Each 3×3 filter is loaded into registers once and then applied across all 36 output positions before the next filter is loaded. Total weight loads: $16 \times 9 = 144$ values, each loaded exactly once.

Ordering B (output-stationary): Place the spatial loops outermost and the filter loops innermost. For every output position, all 16 filters must be loaded, applied, and their partial sums accumulated before advancing to the next position. If the register file cannot hold all 16 filters simultaneously, filters are repeatedly fetched from cache or DRAM. In the worst case, each of the 36 output positions reloads all 144 filter weights, producing $36 \times 144 = 5,184$ weight reads.

Systems insight: Under these two limiting assumptions, Ordering A reduces modeled weight reads by $36 \times$ relative to the worst case for Ordering B. Real kernels also account for activation and partial-sum traffic, register capacity, cache behavior, and vectorization. The example nevertheless shows why two mathematically equivalent loop orders can place very different demands on the memory system.

The combinatorial growth revealed by this analysis poses a practical challenge: explaining how practitioners achieve strong performance despite a large candidate space. Exhaustive enumeration is often impractical, yet production systems routinely find useful schedules for common kernels. The answer lies in a small set of principled dataflow patterns that reduce the search to a manageable set of strategic choices.

11.8 Dataflow Optimization

The mapping strategies from section 11.7 establish *where* computations execute and *where* data resides, but they do not specify **dataflow optimization**: how data flows through processing elements during execution. A systolic array might process a matrix multiplication with weights in local memory, but the order in which weights, inputs, and outputs move through the array directly determines memory bandwidth consumption and energy efficiency. The choice among strategies directly impacts whether an accelerator operates in the compute-bound or memory-bound region

identified by the Roofline analysis—which is why compilers (section 11.9) and runtime systems (section 11.10) must select appropriate dataflow patterns based on workload characteristics.

Three recurring decisions structure the dataflow choices considered here:

1. **Locality:** Weight-stationary, output-stationary, and input-stationary strategies each make different choices about what to cache near compute units, trading off different memory access patterns.
2. **Organization:** Tensor layouts (NHWC vs. NCHW) determine whether memory accesses align with hardware preferences, with performance impacts that can be large when layout conversions or uncoalesced access block the fast path.
3. **Combination:** Kernel fusion and tiling restructure computation to minimize memory traffic, often producing large speedups on low-arithmetic-intensity operations by avoiding intermediate writes and reloads.

These patterns provide a compact vocabulary for reasoning about many dataflow decisions without exhaustive search. The next sections examine each decision in turn, then show how they combine for specific neural network architectures including ResNet-50, GPT-2, and MLPs.

11.8.1 Building blocks of mapping strategies

These three decisions map to four foundational techniques: data movement patterns (weight-stationary, output-stationary, input-stationary), memory-efficient tensor layouts (channels-last vs. channels-first), kernel fusion (combining operations to eliminate intermediate writes), and tiling (partitioning computations into memory-friendly blocks). Together, these building blocks reduce the mapping search space: heuristic and model-driven optimizers can combine them instead of rediscovering the same data-movement choices from scratch.

11.8.1.1 Data movement patterns

While computational mapping determines where and when operations occur, its success depends heavily on how efficiently data is accessed and transferred across the memory hierarchy. Some machine learning workloads are regular but exceed cache capacity; others, such as sparse lookups and routed models, also have irregular access patterns. Both cases make data movement strategy critical to overall system performance.

Even when computational units are mapped efficiently, poor data movement strategies degrade performance by causing frequent memory stalls and leaving hardware resources idle. If data cannot be supplied to processing elements at the required rate, computational units stall, increasing latency, memory traffic, and energy consumption (Y.-H. Chen et al. 2016). Listing 11.15 illustrates how data movement inefficiencies affect the backbone computation of many machine learning models through a typical matrix multiplication operation.

Listing 11.15: Naïve Matrix-Multiplication Dataflow: The loop nest revisits weight rows, input columns, and output accumulators, exposing the reuse that a mapping must capture in local memory.

```
## Matrix multiplication where:
## weights: [512x256] - model parameters
## input:   [256x32]  - batch of activations
## Z:       [512x32]  - output activations

## Computing each output element Z[i,j]:
for i in range(512):
    for j in range(32):
        for k in range(256):
            Z[i, j] += weights[i, k] * input[k, j]
```

This computation reveals several critical dataflow challenges. The first challenge is the number of memory accesses required. For each output Z_{ij} , the computation must fetch an entire row of weights from the weight matrix and a full column of activations from the input matrix. Since the weight matrix contains 512 rows and the input matrix contains 32 columns, this results in repeated memory accesses that place a heavy burden on memory bandwidth.

The second challenge comes from weight reuse. The same weights are applied to multiple inputs, meaning that an ideal mapping strategy should maximize weight locality to avoid redundant memory fetches. Without proper reuse, the accelerator would waste bandwidth loading the same weights multiple times (Chen et al. 2018).

The third challenge involves the accumulation of intermediate results. Since each element in Z_{ij} requires contributions from 256 different weight-input pairs, partial sums must be stored and retrieved before the final value is computed. If these intermediate values are stored inefficiently, the system will require frequent memory accesses, further increasing bandwidth demands.

One way to mitigate these challenges is to use SIMD and SIMT execution models, which allow multiple values to be fetched in parallel. However, even with these optimizations, data movement remains a bottleneck. The primary bottleneck is not retrieval speed alone, but transfer frequency and placement within the memory hierarchy (Han, Liu, et al. 2016).

Because moving data dominates the energy budget that figure 11.9 charts, the single most important goal of an accelerator is to minimize memory access. Dataflow strategies achieve this by maximizing data reuse. The central decision is which data is most valuable to keep local. Accelerators answer that decision by determining which data remains fixed in memory and which data streams dynamically: weight-stationary keeps model parameters local, input-stationary maintains activation data, and output-stationary preserves intermediate results. Each approach trades off different memory access patterns to maximize data reuse and minimize the energy-intensive transfers that constitute the primary bottleneck in AI acceleration.

Weight stationary The weight stationary strategy keeps weights fixed in local memory, while input activations and partial sums are streamed through the system. Weight stationary approaches prove particularly beneficial in CNNs and matrix multiplications, where the same set of weights is applied across multiple inputs. By ensuring weights remain stationary, this method reduces redundant memory fetches, which helps alleviate bandwidth bottlenecks and improves energy efficiency.

A key advantage of weight stationary is that it maximizes weight reuse, reducing the frequency of memory accesses to external storage. Since weight parameters are often shared across multiple computations, keeping them in local memory eliminates unnecessary data movement, lowering the overall energy cost of computation. This makes it particularly effective for architectures where weights represent the dominant memory overhead, such as systolic arrays and custom accelerators designed for machine learning. Listing 11.16 demonstrates how Weight Stationary execution keeps weights fixed in local memory while streaming inputs and accumulating partial sums.

Listing 11.16: Weight Stationary Dataflow: Weights stay resident in local memory while inputs and partial sums stream through, minimizing parameter read traffic; best for CNNs and matrix multiplications with heavy weight reuse.

```
## Weight Stationary Matrix Multiplication
## - Weights remain fixed in local memory
## - Input activations stream through
## - Partial sums accumulate for final output

for weight_block in weights: # Load and keep weights stationary
    load_to_local(weight_block) # Fixed in local storage
    for input_block in inputs: # Stream inputs dynamically
        for output_block in outputs: # Compute results
            output_block += compute(weight_block, input_block)
        # Reuse weights across inputs
```

In weight stationary execution, weights are loaded once into local memory and remain fixed throughout the computation while inputs stream dynamically, reducing redundant memory accesses. Partial sums accumulate efficiently, minimizing unnecessary data movement. Because weights need not be reloaded for each new computation, bandwidth requirements drop significantly, making this dataflow highly effective for workloads with heavy weight reuse patterns such as CNNs and matrix multiplications.

However, while this strategy reduces weight-related memory traffic, it introduces trade-offs in input and output movement. Since inputs must be streamed dynamically while weights remain

fixed, the efficiency of this approach depends on how well input activations can be delivered to the computational units without causing stalls. Partial sums, which represent intermediate results, must also be carefully accumulated to avoid excessive memory traffic. The total performance gain depends on the size of available on-chip memory, as storing larger weight matrices locally can become a constraint in models with millions or billions of parameters.

The weight stationary strategy is well-suited for workloads where weights exhibit high reuse and memory bandwidth is a limiting factor. It is commonly employed in CNNs, systolic arrays, and matrix multiplication kernels, where structured weight reuse leads to measurable performance improvements. However, for models where input or output reuse is more critical, alternative dataflow strategies, such as output stationary or input stationary, may provide better trade-offs.

Output stationary Weight stationary keeps weights local and streams inputs through the system. The dominant cost shifts, however, when the bottleneck is not weight loading but the frequent writes of partial sums. In fully connected layers and transformer attention mechanisms, each output element accumulates contributions from hundreds or thousands of weight-input pairs. Writing those intermediate partial sums to external memory after every accumulation step would create a write-bandwidth bottleneck far more severe than the read overhead that weight stationary addresses. The output stationary strategy inverts the priority: it keeps partial sums fixed in local memory while streaming both weights and input activations through the system, so that each output element is written to external memory only once, after all its contributions have been accumulated (Y.-H. Chen et al. 2016).

Listing 11.17 demonstrates how accumulating partial sums locally minimizes memory writes and enhances efficiency during matrix multiplication. In this implementation, the accumulator buffer stays in local registers or scratchpad throughout the inner loop; weights and inputs stream in, contribute to the running sum, and are discarded. The final result is written out only once per output element, eliminating the repeated write traffic that would otherwise dominate bandwidth.

This approach aligns naturally with systolic arrays, where computation progresses through a grid of processing elements and partial sums can flow along one axis without leaving the chip. The trade-off is that both weights and activations must now be streamed dynamically, so the system must sustain high read bandwidth for two data streams simultaneously. Parallel implementations also require careful synchronization when multiple PEs contribute to the same output element. Output stationary is therefore most effective for workloads where accumulation dominates, such as fully connected layers and attention mechanisms, but less suitable when input reuse is the critical bottleneck.

Listing 11.17: Output Stationary Dataflow: Partial sums stay resident in local memory while weights and inputs stream through, so each output is written out only once, minimizing accumulation write traffic; best for fully connected layers and attention.

```
## - Partial sums remain in local memory
## - Weights and input activations stream through dynamically
## - Final outputs are written only once

for output_block in outputs: # Keep partial sums stationary
    accumulator = 0 # Initialize accumulation buffer
    for weight_block, input_block in zip(weights, inputs):
        accumulator += compute(weight_block, input_block)
        # Accumulate partial sums
    store_output(accumulator) # Single write to memory
```

Input stationary The two strategies examined so far each fix a different operand in local memory: weight stationary fixes weights to reduce read bandwidth for parameters, and output stationary fixes partial sums to reduce write bandwidth for accumulations. The third strategy completes the picture by fixing the remaining operand: input activations. Within a transformer layer, one activation can feed multiple projections or attention heads; the next layer receives a newly computed representation rather than the same token activation. When activation reuse is the dominant memory cost, keeping inputs stationary and streaming weights through the system can reduce energy and bandwidth. Listing 11.18 illustrates this approach.

Listing 11.18: Input Stationary Dataflow: Input activations stay resident in local memory while weights stream through, reducing activation read traffic when one activation tile feeds several computations.

```
## - Input activations remain in local memory
## - Weights stream through dynamically
## - Partial sums accumulate and are written out

for input_block in inputs: # Keep input activations stationary
    load_to_local(input_block) # Fixed in local storage
    for weight_block in weights: # Stream weights dynamically
        for output_block in outputs: # Compute results
            output_block += compute(weight_block, input_block)
        # Reuse inputs across weights
```

Here, input activations are loaded once and held fixed while weights stream through. Partial sums accumulate and are eventually written out, but unlike output stationary, the accumulation buffer is not the primary beneficiary of locality; instead, the input data is.

The trade-off mirrors the other two strategies: weights must now be streamed dynamically, so the system needs sustained read bandwidth for the weight stream, and partial sums require buffering before write-back. Input stationary is most effective where an activation tile feeds multiple projections, heads, gates, or other computations before eviction.

Taken together, the three dataflow strategies provide a decision rule rather than a hierarchy of quality. Weight stationary minimizes read traffic for parameters and suits CNNs with small, heavily reused filters. Output stationary minimizes write traffic for accumulations and suits fully connected layers with high fan-in. Input stationary minimizes read traffic for activations and suits transformers and batch processing with high activation reuse. No single strategy dominates; the optimal choice depends on which data element has the highest reuse ratio relative to its size, a determination that the compiler and hardware designer must make based on the specific workload and memory hierarchy. Convolution-specific designs add a fourth label, row-stationary (as in Eyeriss), but it is not a separate primitive: it is a hybrid that keeps rows of inputs and partial sums local when that mapping yields higher reuse than fixing any single operand (section 11.5.6.2).

11.8.1.2 Memory-efficient tensor layouts

The dataflow strategies in section 11.8.1.1 determine which data stays close to compute; tensor layouts determine whether that data can be accessed efficiently once it arrives. A perfectly chosen weight-stationary dataflow still suffers if weights are stored in a format that causes scattered memory accesses. Tensor layout is therefore a kernel contract: the physical arrangement of multidimensional data must match the access pattern expected by the selected hardware path, or the accelerator pays in memory stalls, inefficient cache usage, and increased data movement.

In AI accelerators, tensor layout optimization is particularly important because data is frequently accessed in patterns dictated by the underlying hardware architecture. Choosing the right layout ensures that memory accesses align with hardware-friendly access patterns, minimizing overhead from costly memory transactions (NVIDIA Corporation 2021b).

While developers can sometimes manually specify tensor layouts, the choice is often determined automatically by machine learning frameworks such as TensorFlow, PyTorch, and JAX, by compilers, or by AI accelerator runtimes. Low-level optimization tools such as cuDNN (for NVIDIA GPUs), XLA (for TensorFlow graphs), and MLIR-based compiler stacks may impose or transform tensor layouts as they lower operations to backend-specific kernels (NVIDIA Corporation 2021b; Google 2025; Lattner et al. 2020). In high-level frameworks, layout transformations are typically applied transparently, but developers working with custom kernels or low-level libraries such as CUDA, Metal, or OpenCL may have direct control over tensor format selection.

For example, PyTorch exposes view operations such as `tensor.permute()` and uses `tensor.contiguous()` to materialize contiguous storage when a downstream kernel requires it (PyTorch Contributors 2026b). TensorFlow recommends channels-last (NHWC) for convolutional models on NVIDIA GPUs and warns that NCHW can introduce automatic transpose overhead (TensorFlow Developers 2024b). Hardware-aware libraries such as cuDNN for GPUs and oneDNN for CPUs use specific memory layouts to maximize cache locality and SIMD efficiency. The practical rule is to treat layout as part

of the selected backend path: the fastest tensor format is the one that avoids conversion overhead and makes the kernel's memory accesses contiguous.

Channels-last layout In an NHWC channels-last layout, the channel index varies fastest in a contiguous tensor, so the channel values for each pixel are adjacent. NHWC describes logical dimension order, not whether the underlying linear storage is row-major; both NHWC and NCHW tensors can use row-major linearization of their respective dimensions.

To understand channels-last layout, consider a single RGB image represented as a tensor of shape (Height, Width, Channels). If the image has a size of 3×3 pixels with 3 channels (RGB), the corresponding tensor is structured as (3, 3, 3). The values are stored in memory as follows:

$$I(0,0,0), I(0,0,1), I(0,0,2), I(0,1,0), I(0,1,1), \\ I(0,1,2), I(0,2,0), I(0,2,1), I(0,2,2), \dots$$

The channels for each pixel are adjacent, and pixels advance across width before height. This ordering can support sequential access when a kernel vectorizes over adjacent channel values.

The benefit of channels-last storage depends on the operator and backend. CPU and accelerator kernels that vectorize over channels can use the contiguous channel dimension, while other kernels may prefer channels-first or blocked internal layouts.

However, layout choice becomes subtle for convolutions because logical dimension order and physical memory format are not the same thing. A tensor may be described as NHWC or NCHW, but the backend ultimately cares whether the memory addresses consumed by a kernel are contiguous, aligned, and coalesced for the specific operator and precision mode.

TensorFlow commonly uses NHWC conventions, while PyTorch commonly exposes NCHW tensors with a separate channels-last memory format option. When targeting GPUs, frameworks and libraries may insert, propagate, or internally perform layout transformations to match the fastest kernel path.

Channels-first layout In a contiguous NCHW channels-first layout, the width index varies fastest and all spatial values for a channel are stored together before the next channel. GPUs process data in parallel across threads, and when threads access consecutive memory addresses, the hardware can combine these requests into a single efficient transaction (memory coalescing). Historically, many GPU convolution paths used NCHW effectively, while modern Tensor Core convolution and fusion paths often prefer NHWC or channels-last physical layouts because those layouts align better with vectorized tensor-core kernels.

To understand channels-first layout, consider the same RGB image tensor with dimensions reordered as (Channels, Height, Width) = (3, 3, 3). Retaining $I(h, w, c)$ as semantic coordinate notation, the data are stored in channels-first order as follows:

$$I(0,0,0), I(0,1,0), I(0,2,0), I(1,0,0), I(1,1,0), I(1,2,0), \dots, \\ I(0,0,1), I(0,1,1), I(0,2,1), \dots, I(0,0,2), I(0,1,2), I(0,2,2), \dots$$

In this format, all red channel values for the entire image are stored first, followed by all green values, and then all blue values. This ordering can allow some hardware accelerators to efficiently load and process data across channels in parallel, which is important for convolution operations and SIMD execution models (Chetlur et al. 2014).

The advantage of a given layout becomes clear only relative to a specific backend. Convolutional layers process images by applying a shared set of filters across all channels. Depending on the kernel implementation, NCHW, NHWC, or a blocked internal layout may minimize scattered memory fetches, reduce memory latency, and improve data locality for the lowered matrix multiplications that implement convolution.

Because GPUs and TPUs rely on **memory coalescing**,³⁴ a technique in which consecutive threads fetch contiguous memory addresses, the best layout is the one that makes the kernel's actual thread-access pattern contiguous. For example, in NVIDIA GPU convolution paths, cuDNN may use or internally convert to NHWC/channels-last for Tensor Core kernels, while other kernels may still perform well with NCHW. The rule is backend- and operator-dependent rather than a universal CPU=NHWC, GPU=NCHW split.

³⁴ | **Memory coalescing:**

The GPU hardware mechanism that fuses memory requests from threads in a warp into a single transaction when those threads access contiguous memory. Tensor layout affects coalescing, but the sign of the effect depends on the kernel: NCHW can be efficient for some convolution implementations, while NHWC/channels-last is often preferred for modern Tensor Core convolution and fusion paths. Poor layout choices can still create multi-fold performance gaps, but the correct fix is to match the layout to the backend rather than memorize one universal format.

Despite its advantages for some accelerator kernels, channels-first layout can introduce inefficiencies in kernels optimized for channels-last access. The most efficient layout depends on the operation, framework convention, and library implementation.

Modern AI frameworks and compilers often transform tensor layouts dynamically depending on the execution environment, but this is not guaranteed for every model or operation.³⁵ TensorFlow, XLA, cuDNN, and TensorRT may insert or choose layout conversions internally; PyTorch exposes explicit channels-last conversion and propagation paths. Developers still need to profile layout choices when convolution performance is material.

Comparing channels-last and channels-first layouts Both channels-last (NHWC) and channels-first (NCHW) layouts serve distinct purposes in machine learning workloads, with their efficiency largely determined by the hardware architecture, memory access patterns, and computational requirements. The choice of layout directly influences cache utilization, memory bandwidth efficiency, and processing throughput. Table 11.18 contrasts the performance trade-offs and hardware compatibility between these two approaches.

Table 11.18: Data Layout Strategies: Channels-last (NHWC) and channels-first (NCHW) layouts optimize memory access patterns for different backend kernels. NHWC often suits CPU vectorization and modern Tensor Core convolution/fusion paths, while NCHW remains common in PyTorch model code and many GPU kernels. Choosing the appropriate layout directly impacts performance by maximizing cache utilization and memory bandwidth efficiency.

Feature	Channels-Last (NHWC)	Channels-First (NCHW)
Memory storage order	Pixels are stored row-by-row, channel interleaved	All values for a given channel are stored together first
Best for	CPU loops, element-wise operations, many channels-last kernels	Many legacy GPU convolution paths and channel-first model code
Cache efficiency	High cache locality for sequential row/channel-last access	Can improve coalescing for channel-first kernels
Convolution performance	Often preferred by modern Tensor Core convolution/fusion paths	Efficient for many cuDNN and framework kernels
Memory fetching	Good when kernels vectorize over adjacent channel data	Good when kernels process channel-major tiles
Default in frameworks	Common TensorFlow convention; PyTorch channels-last option	Common PyTorch tensor shape convention

The decision to use channels-last (NHWC) or channels-first (NCHW) layouts is not always made manually by developers. Instead, machine learning frameworks and AI compilers often determine the optimal layout based on the target hardware and operation type.

In practice, modern AI compilers such as TensorFlow’s XLA, cuDNN, TensorRT, and PyTorch compilation paths may perform layout transformations or propagate layout metadata. The result can be high throughput without manual tensor rewrites, but performance-sensitive deployments should still profile both layout and conversion overhead on the target hardware.

11.8.1.3 Kernel fusion

One of the most impactful optimization techniques in AI acceleration involves reducing the overhead of intermediate data movement between operations. Kernel fusion³⁶ transforms multiple separate computations into unified operations, dramatically improving memory efficiency and execution performance. The memory bottlenecks created by intermediate writes motivate kernel fusion, which eliminates these inefficiencies.

Intermediate memory write AI model performance is often constrained by memory bandwidth and intermediate memory writes rather than pure arithmetic operations. Every time an operation produces an intermediate result that must be written to memory and later read back, execution stalls from the data movement overhead.

Kernel fusion represents the critical bridge between the software optimization techniques introduced in section 10.5.1.5 and the memory bandwidth constraints analyzed in section 11.5.1. Many AI workloads introduce unnecessary intermediate memory writes, increasing memory bandwidth consumption and reducing execution efficiency (NVIDIA Corporation 2017).

35 | **NHWC vs. NCHW:** NHWC lists dimensions as batch, height, width, channel; NCHW lists batch, channel, height, width. Physical memory format determines whether adjacent threads read adjacent addresses, so the performance effect is backend-specific. Layout-to-hardware mismatch is not a micro-optimization; it can create multi-fold performance gaps, especially when a conversion prevents Tensor Core convolution or fusion paths from being used.

36 | **Kernel fusion:** Separate GPU kernels expose intermediate results through device memory, although caches can absorb some traffic. A fused kernel can keep intermediates in registers or local storage and avoid corresponding external write/read cycles. The resulting speedup depends on cache behavior, occupancy, launch overhead, and whether memory traffic was the bottleneck; it is not necessarily proportional to the byte reduction.

Listing 11.19 reveals how each operation becomes a separate kernel in a naïve execution model, forcing intermediate results to be written to memory and then read back for the next operation. The first two operations produce intermediate tensors that a naïve multi-kernel execution exposes through device memory for the next operation. On large tensors, this movement can outweigh the arithmetic cost (Jia et al. 2019). Table 11.19 illustrates the live-tensor footprint when buffers are not reused.

Listing 11.19: Naïve Execution: Each step writes intermediate results to memory before processing the next, leading to increased bandwidth usage and reduced efficiency.

```
import torch
import torch.nn.functional as F

## Input tensor
X = torch.randn(1024, 1024).cuda()
running_mean = torch.zeros(1024, device=X.device)
running_var = torch.ones(1024, device=X.device)

## Step-by-step execution (naïve approach)
X1 = torch.relu(X) # Intermediate tensor stored in memory
X2 = F.batch_norm(
    X1, running_mean, running_var, training=False
) # Frozen inference BatchNorm
Y = 2.0 * X2 + 1.0 # Final result
```

Kernel fusion for memory efficiency The two intermediate tensors consume memory capacity when they remain live and create external traffic when separate kernels materialize them. Kernel fusion can eliminate those intermediate writes and reloads (Jia et al. 2019) by propagating values through registers or local memory within one kernel.

Table 11.19: Intermediate Tensor Storage: If all four named tensors remain live in this naïve example, they occupy 16.8 MB; an optimizing runtime may reuse buffers when earlier values are no longer needed. Fusion primarily guarantees removal of the external write and read for each eliminated intermediate rather than a universal reduction in peak allocated memory.

Tensor	Size (MB) for 1024×1024 Tensor
X	4.2 MB
X'	4.2 MB
X''	4.2 MB
Y	4.2 MB
Total Memory	16.8 MB

A machine learning inference sequence might apply ReLU, frozen batch normalization, and then an affine scaling with scale α and offset β . In a naïve implementation, each separate kernel generates an intermediate tensor that is written to memory and read back by the next kernel:

$$\begin{aligned} X' &= \text{ReLU}(X) \\ X'' &= \text{BatchNorm}(X') \\ Y &= \alpha \cdot X'' + \beta \end{aligned}$$

With kernel fusion, these operations are combined into a single computation step, allowing the entire transformation to occur without generating unnecessary intermediate tensors:

$$Y = \alpha \cdot \text{BatchNorm}(\text{ReLU}(X)) + \beta$$

Table 11.20 highlights the impact of operation fusion on memory efficiency. By keeping intermediate results in registers or local memory rather than writing them to main memory, fusion reduces external traffic. For three separate pointwise inference kernels, the idealized naïve path reads and writes one tensor per operation, whereas the fused path reads the input once and writes the output once.

Table 11.20: Operation Fusion Benefits: In this idealized inference sequence, fusion reduces external tensor traffic from 25.2 MB to 8.4 MB, a $3\times$ reduction. Actual traffic depends on cache behavior and the compiler's buffer reuse.

Execution Model	External tensor payloads	Traffic (MB)
Naïve execution	3 reads + 3 writes	25.2 MB
Fused execution	1 read + 1 write	8.4 MB

Performance benefits and constraints Kernel fusion brings several key advantages that enhance memory efficiency and computation throughput. By reducing memory accesses, fused kernels ensure that intermediate values stay within registers instead of being repeatedly written to and read from memory. This significantly lowers memory traffic, which is one of the primary bottlenecks in machine learning workloads. GPUs and TPUs, in particular, benefit from kernel fusion because high-bandwidth memory is a scarce resource, and reducing memory transactions leads to better utilization of compute units (NVIDIA Corporation 2020a).

However, not all operations can be fused arbitrarily. Element-wise operations and frozen inference normalization are strong candidates because their computations do not require batch-wide reductions. Training-mode batch normalization computes batch statistics and is not purely element-wise. Matrix multiplications and convolutions constrain fusion because they involve reductions and large data movement; they are often fused with element-wise epilogues such as bias or activation, but cannot be freely fused with unrelated global operations.

Another major consideration is register pressure. Fusing multiple operations means all temporary values must be kept in registers rather than memory. While this eliminates redundant memory writes, it also increases register demand. If a fused kernel exceeds the available registers per thread, the compiler can spill excess values to per-thread local memory, which is backed by device memory and may be cached, introducing additional latency and potentially negating the benefits of fusion. On GPUs, where thread occupancy (the number of threads that can run in parallel) is limited by available registers, excessive fusion can reduce parallelism, leading to diminishing returns.

Different AI accelerators and compilers handle fusion in distinct ways. NVIDIA GPUs, for example, favor warp-level parallelism, where element-wise fusion is straightforward (NVIDIA Corporation 2020a). TPUs, on the other hand, prioritize systolic array execution for dense matrix operations (Jouppi et al. 2017). Compiler and inference stacks such as TVM, XLA, TensorRT, and MLIR apply graph rewrites, lowering passes, or engine-building heuristics to balance memory savings against execution constraints (Chen et al. 2018; Google 2025; NVIDIA 2024c; Lattner et al. 2020).

Despite its advantages, fusion is not always beneficial. Some AI frameworks allow developers to disable fusion selectively, especially when debugging performance issues or making frequent model modifications. The decision to fuse operations must consider trade-offs between memory efficiency, register usage, and hardware execution constraints to ensure that fusion leads to tangible performance improvements.



Checkpoint 11.3: Data movement and kernel fusion

The dataflow building blocks are now in place: data locality, tensor layout, and kernel fusion. These fusion decisions are ultimately about data locality, which ties together the chapter's core data movement strategies:

- ☐ **Locality choice:** given weight-stationary, output-stationary, and input-stationary dataflows, which tensor does each hold near the compute units, and how does that choice shift where the memory traffic lands?
- ☐ **Layout match:** a kernel runs on a backend expecting NHWC but receives an NCHW tensor. Predict what happens to its memory access pattern and to delivered throughput.
- ☐ **Fusion eligibility:** for a Conv2D, a BatchNorm, and a ReLU executed in sequence, decide which pairs are worth fusing and what determines whether fusing them pays off.
- ☐ **Remaining gap:** locality, layout, and fusion are all chosen well, yet the operation still stalls on memory. What is the missing transformation, and why do the other three not address it?

11.8.1.4 Memory-efficient tiling strategies

While modern AI accelerators offer high computational throughput, their performance is often limited by memory bandwidth rather than raw processing power. If data cannot be supplied to processing units fast enough, execution stalls occur, leading to wasted cycles and inefficient hardware utilization.

Tiling³⁷ mitigates this issue by restructuring computations into smaller, memory-friendly sub-problems. When memory bandwidth cannot be increased directly, systems must reduce trips to main memory. Instead of processing entire matrices or tensors at once, which leads to excessive memory traffic, tiling partitions computations into smaller blocks (tiles) that fit within fast local memory (for example, caches, shared memory, or registers) (Lam et al. 1991).

Matrix multiplication, widely used in AI models, demonstrates inefficient memory access when implemented naively. Listing 11.20 shows how, without tiling, repeated memory accesses for the same data lead to unnecessary bandwidth consumption.

Listing 11.20: Naïve Matrix Multiplication: Direct implementation makes $\mathcal{O}(N^3)$ array references for $N \times N$ matrices. Caches can serve some references, but the loop order does not deliberately block reuse for the memory hierarchy.

```
for i in range(N):
    for j in range(N):
        for k in range(N):
            C[i, j] += A[i, k] * B[k, j]  # Repeatedly fetching
            # A[i, k] and B[k, j]
```

The loop repeatedly references elements of **A** and **B**. Some references can hit in cache, but large matrices and an unfavorable traversal order can still cause excessive lower-tier traffic.

Tiling addresses this problem by ensuring that smaller portions of matrices are loaded into fast memory, reused efficiently, and only written back to main memory when necessary. This technique is especially important in AI accelerators, where memory accesses dominate execution time. Figure 11.14 labels the product $\mathbf{C} = \mathbf{A}\mathbf{B}$ by its three dimensions: M rows and K columns in **A**, K rows and N columns in **B**, and $M \times N$ in the output **C**. Each highlighted tile is the working set that fits in fast memory at one moment: a green $M_{\text{tile}} \times K_{\text{tile}}$ row band of **A** multiplies a pink $K_{\text{tile}} \times N_{\text{tile}}$ column band of **B** to accumulate one blue $M_{\text{tile}} \times N_{\text{tile}}$ output block ($\text{Block}_{m,n}$) of **C**. Processing all computations for one tile before moving to the next avoids repeatedly paying the DRAM access penalty.

Tiling fundamentals Tiling divides a computation into smaller tiles that fit within available fast memory rather than operating on an entire data structure at once. This structure maximizes data reuse, reducing redundant memory accesses and improving efficiency.

Consider matrix multiplication, a key operation in machine learning workloads. The operation computes $\mathbf{C} = \mathbf{A} \times \mathbf{B}$ where each element $C_{ij} = \sum_k A_{ik} \times B_{kj}$. Listing 11.20 exposes the core problem: every iteration of the innermost loop fetches elements from matrices **A** and **B** from memory, performs a multiplication, and updates matrix **C**. Because matrices are large, the processor repeatedly reloads the same values from memory, even though they were just used in previous computations.

This data movement overhead is expensive: DRAM access has much higher latency and energy cost than access to on-chip cache or registers (Horowitz 2014; Sze et al. 2017). The solution is tiling.

Performance benefits of tiling Instead of computing one element at a time and constantly moving data in and out of slow memory, tiling processes submatrices (tiles) at a time, keeping frequently used values in fast memory. The idea is to divide the matrices into smaller blocks that fit within the processor’s cache or shared memory, ensuring that once a block is loaded, it is reused multiple times before moving to the next one. Listing 11.21 demonstrates cache-friendly loop blocking: the loop bounds partition the matrices into tiles, and the hardware cache hierarchy keeps recently used tile data close to the compute units when access order has locality.

³⁷ | **Tiling (loop blocking):** This restructuring partitions a computation into blocks sized for a fast local tier. A naive loop makes repeated references to the same elements; caches may capture some, while blocking deliberately concentrates reuse before eviction (Lam et al. 1991). This reduction in lower-tier traffic is a primary source of the gap between naive matrix multiplication and optimized general matrix multiplication (GEMM) routines.




Figure 11.14: Matrix Tiling Mechanics: Partitioning large matrix multiplication ($C = AB$) into $M_{\text{tile}} \times K_{\text{tile}}$ and $K_{\text{tile}} \times N_{\text{tile}}$ blocks aligns memory footprints with fast on-chip SRAM/registers. Maximizing register reuse within each tile amortizes slow off-chip DRAM accesses and keeps compute engines operating at peak throughput.

Listing 11.21: Cache-Blocked Matrix Multiplication: This high-level loop blocking approach divides matrices into smaller index ranges so hardware caches can reuse data within each tile, improving computational efficiency without explicit scratchpad loads.

```
TILE_SIZE = 32 # Choose a tile size based on hardware constraints

# Cache blocking: partition data via loop bounds.
# Loads are implicit through the hardware cache hierarchy.
for i in range(0, N, TILE_SIZE):
    for j in range(0, N, TILE_SIZE):
        for k in range(0, N, TILE_SIZE):
            # Each tile computed independently
            for ii in range(i, min(i + TILE_SIZE, N)):
                for jj in range(j, min(j + TILE_SIZE, N)):
                    for kk in range(k, min(k + TILE_SIZE, N)):
                        C[ii, jj] += A[ii, kk] * B[kk, jj]
```

This restructuring significantly improves performance through three reinforcing effects. Memory reuse improves because the approach visits a small tile repeatedly while it is likely to remain in cache before moving on to the next tile, rather than fetching elements from **A** and **B** repeatedly from slow memory, which minimizes redundant memory accesses. Memory bandwidth usage drops as a direct consequence: since each tile is used multiple times before being evicted, most required data is available in L1/L2 cache or shared memory rather than DRAM, so traffic falls and execution speeds up. Compute efficiency rises in turn, because processors spend less time waiting for data and more time performing useful work; in architectures like GPUs and TPUs, where thousands of parallel processing units operate simultaneously, tiling keeps data read and processed in a structured manner that avoids unnecessary stalls.

This technique is particularly effective in AI accelerators, where machine learning workloads consist of large matrix multiplications and tensor transformations. Without tiling, these workloads

quickly become memory bound, meaning performance is constrained by how fast data can be retrieved rather than by the raw computational power of the processor.

Tiling methods While the general principle of tiling remains the same, which involves partitioning large computations into smaller subproblems to improve memory reuse, there are different ways to apply tiling based on the structure of the computation and hardware constraints. The two primary tiling strategies are spatial tiling and temporal tiling. These strategies optimize different aspects of computation and memory access, and in practice, they are often combined to achieve the best performance.

Spatial tiling partitions data structures into smaller blocks that fit within fast memory. The tiled matrix multiplication in listing 11.21 demonstrates cache-friendly loop blocking: the code does not issue explicit scratchpad loads, but the tile-shaped access pattern lets hardware caches reuse nearby values before they are evicted. This strategy is particularly beneficial for large tensors that exceed fast memory capacity—by breaking computations into smaller tiles, data movement between memory levels is minimized, keeping operations localized within cache hierarchies.

Temporal tiling complements spatial tiling by explicitly staging data in shared memory or registers and reorganizing the computation order around that staged data. Many ML workloads access the same data repeatedly across iterations—without temporal tiling, this results in redundant memory fetches. Temporal tiling restructures the computation to ensure that frequently used data stays in fast memory for as long as possible before the next computation begins.

A classic example where temporal tiling is beneficial is convolutional operations, where the same set of weights is applied to multiple input regions. Without loop blocking, these weights might be loaded from memory multiple times for each computation. With temporal tiling, the computation is reordered so that the weights remain in fast memory across multiple inputs, reducing unnecessary memory fetches and improving overall efficiency. Listing 11.22 illustrates explicit tile staging: the code loads blocks of **A** and **B** into temporary fast storage, then reuses them across multiple inner-loop operations.

Listing 11.22: Explicit Tile Staging Pseudocode: A backend maps the temporary tile objects to fast storage and reuses them across inner-loop operations.

```
# Explicit tile staging: load data into fast
# temporary storage before the inner loops.
for i in range(0, N, TILE_SIZE):
    for j in range(0, N, TILE_SIZE):
        for k in range(0, N, TILE_SIZE):
            # Pseudocode: map these tile objects to fast memory
            A_tile = A[i : i + TILE_SIZE, k : k + TILE_SIZE]
            B_tile = B[k : k + TILE_SIZE, j : j + TILE_SIZE]

            # Reuse loaded tiles for all inner iterations
            for ii in range(A_tile.shape[0]):
                for jj in range(B_tile.shape[1]):
                    for kk in range(A_tile.shape[1]):
                        C[i + ii, j + jj] += (
                            A_tile[ii, kk] * B_tile[kk, jj]
                        )
```

Explicit tile staging improves performance when the backend places the temporary tiles in fast memory and reuses them before eviction. The pseudocode includes shortened boundary tiles when *N* is not divisible by *TILE_SIZE*.

This technique is particularly useful in workloads where certain values are used repeatedly, such as convolutions, recurrent neural networks (RNNs), and self-attention mechanisms in transformers. By applying loop blocking, AI accelerators can significantly reduce memory stalls and improve execution throughput.

Tiling challenges and trade-offs Tiling improves performance only when the tile matches the locality budget of the hardware. If the tile is too small, memory fetches still dominate execution time because reuse is too limited. If the tile is too large, it exceeds fast memory and causes cache thrashing

or scratchpad spills. Selecting the right tile size therefore directly determines computational efficiency and memory bandwidth usage.

The tile choice also controls load balance. In architectures such as GPUs and TPUs, computations execute in parallel across thousands of processing units. If tiles are not evenly distributed, some units remain idle while others are overloaded, leading to suboptimal utilization of computational resources. Effective tile scheduling keeps parallel execution balanced and efficient.

Data movement remains the limiting cost even after tiling. Although tiling reduces the number of slow memory accesses, transferring tiles between hierarchy levels still incurs latency and energy cost, especially when data falls from cache or scratchpad back to DRAM. Efficient memory prefetching and scheduling strategies minimize this residual movement and ensure that data is available when needed.

Hybrid tiling combines spatial and temporal strategies when neither dimension alone captures the workload's reuse pattern. Some AI accelerators use spatial tiling for matrix multiplications while employing temporal tiling for weight reuse in convolutional layers, dynamically adjusting tile sizes or reordering computations based on real-time execution conditions.

Register blocking, double buffering, and hierarchical tiling extend the same locality principle at smaller and larger memory tiers. AI compilers and runtime systems such as TensorFlow XLA, TVM, and MLIR automatically select these tiling strategies based on hardware constraints, enabling fine-tuned performance optimization without manual intervention. Table 11.21 provides a comparative overview of spatial, temporal, and hybrid tiling approaches, highlighting their respective benefits and trade-offs.

When machine learning models grow in size and complexity, tiling remains a critical tool for improving hardware efficiency, ensuring that AI accelerators operate near their practical potential. While manual tiling strategies can provide substantial benefits, compilers and hardware-aware optimization techniques further enhance performance by automatically selecting effective tiling strategies for a given workload.

Table 11.21: Tiling Strategies: Spatial, temporal, and hybrid tiling optimize memory access patterns for improved performance. Spatial tiling maximizes data reuse within fast memory, temporal tiling exploits loop structure for reduced accesses, and hybrid tiling combines both approaches. AI compilers and runtime systems use these techniques to automatically optimize model execution on diverse hardware.

Aspect	Spatial Tiling (Data Tiling)	Temporal Tiling (Loop Blocking)	Hybrid tiling
Primary Goal	Reduce memory accesses by keeping data in fast memory longer	Increase data reuse across loop iterations	Adapt dynamically to workload constraints
Optimization Focus	Partitioning data structures into smaller, memory-friendly blocks	Reordering computations to maximize reuse before eviction	Balancing spatial and temporal reuse strategies
Memory usage	Improves cache locality and reduces DRAM access	Keeps frequently used data in fast memory for multiple iterations	Minimizes data movement while ensuring high reuse
Common use cases	Matrix multiplications, CNNs, self-attention in transformers	Convolutions, recurrent neural networks (RNNs), iterative computations	AI accelerators with hierarchical memory, mixed workloads
Performance gains	Reduced memory bandwidth requirements, better cache utilization	Lower memory fetch latency, improved data locality	Maximized efficiency across multiple hardware types
Challenges	Requires careful tile size selection, inefficient for workloads with minimal spatial reuse	Can increase register pressure, requires loop restructuring	Complexity in tuning tile size and execution order dynamically
Best when	Data is large and needs to be partitioned for efficient processing	The same data is accessed multiple times across iterations	Both data partitioning and iteration-based reuse are important

11.8.2 Applying mapping strategies to neural networks

While these foundational mapping techniques apply broadly, their effectiveness varies based on the computational structure, data access patterns, and parallelization opportunities of different neural network architectures. Each architecture imposes distinct constraints on data movement,

memory hierarchy, and computation scheduling, requiring tailored mapping strategies to optimize performance.

A structured approach to mapping is required to address the combinatorial explosion of choices that arise when assigning computations to AI accelerators. Rather than treating each model as a separate optimization problem, the same principles apply across different architectures; only their priority shifts based on workload characteristics. The goal is to systematically select and apply mapping strategies that maximize efficiency for different types of machine learning models.

These principles apply to three representative AI workloads, each characterized by distinct computational demands. CNNs benefit from spatial data reuse, making weight-stationary execution and tiling especially effective. Transformer behavior depends on phase and shape: training and prefill matrix multiplications can be compute bound, while low-batch decoding often depends on weight and KV-cache bandwidth. MLPs involve substantial matrix multiplication and benefit from structured tiling, optimized weight layouts, and memory-aware execution.

Despite their differences, each of these models follows a common set of mapping principles, with variations in how optimizations are prioritized. Table 11.22 summarizes the suitability of different optimization strategies for CNNs, transformers, and MLPs.

11.8.2.1 Convolutional neural networks (ResNet-50)

For ResNet-50 and similar CNNs, a key mapping opportunity is spatial weight reuse. A small filter is applied across many spatial locations, so its weights participate in many multiply-accumulate operations. Weight-stationary mappings can capture that reuse by retaining filter tiles locally, while row-stationary or output-stationary mappings may better balance activation and partial-sum traffic on other architectures. Favorable convolution shapes can achieve high arithmetic intensity, which makes them well suited to matrix and systolic execution paths. The model family does not determine one universal dataflow: layer shape, batch size, local-storage capacity, and the target’s supported kernels decide which reuse pattern is worth retaining.

Table 11.22: Architecture-Specific Mapping Strategies: Each neural network architecture benefits from different optimization priorities based on its computational and memory characteristics.

Optimization technique	CNNs	Transformers	MLPs	Rationale
Dataflow strategy	Weight, row, or output stationary	Phase- and kernel-dependent	Weight or output stationary	CNNs expose filter and spatial reuse; transformer attention and MLP blocks impose different traffic; dense MLPs reuse weights across a batch and accumulate outputs.
Memory-aware tensor layouts	Backend-dependent (often NCHW for cuDNN convolutions, channels-last on Tensor Cores)	Backend-dependent (row-major activations typical)	Row-major typical	Layout choice depends on backend kernel path and precision mode (see section 11.8.1.2); the entries name common defaults, not universal prescriptions.
Kernel Fusion	Convolution + Activation	Fused Attention	GEMM Fusion	CNNs optimize convolution+activation fusion; Transformers fuse attention mechanisms; MLPs benefit from fused matrix multiplications.
Tiling for Memory Efficiency	Spatial Tiling	Temporal Tiling	Blocked Tiling	CNNs tile along spatial dimensions; Transformers use loop blocking to improve sequence memory efficiency; MLPs use blocked tiling for large matrix multiplications.

This spatial regularity also enables aggressive fusion and tiling. For inference with frozen Batch-Norm statistics, compilers can fold or fuse convolution, normalization, and activation to avoid unnecessary intermediate writes. Training-mode BatchNorm requires reductions over batch statistics and does not have the same element-wise fusion contract. Spatial tiling partitions the feature map into subregions sized to fit within on-chip SRAM, so the fused kernel processes each tile from fast memory before moving to the next.

11.8.2.2 Transformer architectures (GPT-2/Llama)

Where CNNs are defined by weight reuse, transformers are defined by the memory pressure of the key-value (KV) cache. During attention computation, every query vector must access stored key and

value pairs across the entire sequence length. As sequences grow, the full KV cache usually lives in HBM or device DRAM, while attention kernels tile active blocks through SRAM, registers, or shared memory. This access pattern motivates activation stationary execution: keep the currently used KV tiles close to compute while streaming queries through them, rather than repeatedly materializing large attention intermediates in external memory.

The memory traffic created by standard attention implementations explains why fused attention kernels, such as FlashAttention (T. Dao et al. 2022), can deliver large performance gains: by fusing the query-key dot product, softmax normalization, and value-weighted summation into a single kernel that tiles along the sequence dimension, these implementations avoid materializing the full attention matrix in main memory. This temporal tiling approach processes sequence blocks that fit within on-chip SRAM, substantially reducing HBM traffic while preserving the $\mathcal{O}(S^2)$ attention computation. Transformer mapping depends on the phase and shape: training and prefill GEMMs can be compute-bound, while low-batch autoregressive decode is commonly constrained by weight and KV-cache bandwidth.

11.8.2.3 Multilayer perceptrons and DLRM

MLPs present the most straightforward mapping problem because their computation reduces largely to dense GEMM.³⁸ Each fully connected layer multiplies an activation matrix by a weight matrix. The weight matrix is fixed across samples, so batching reuses each loaded weight across more activations and raises arithmetic intensity. A batch size of one provides little cross-sample weight reuse and often underutilizes wide matrix engines, while larger batches can move execution toward the compute-bound regime. Section C.1.4 derives the scaling that governs this sensitivity.

Because MLP layers are typically followed by activation functions and bias additions, GEMM fusion combines these steps into a single kernel, avoiding intermediate memory writes. Blocked tiling partitions the large matrix multiplications into sub-blocks sized for the accelerator’s shared memory, ensuring high cache utilization throughout computation. The simplicity of the MLP mapping, dominated by a single primitive with predictable access patterns, is precisely why hardware vendors optimize GEMM libraries so aggressively: gains in GEMM performance translate directly to MLP throughput. DLRM is not purely an MLP workload, however. Its large sparse embedding tables and feature-interaction stage can dominate capacity and memory traffic, so the dense mapping described here applies to its bottom and top MLP towers rather than to the entire model (Naumov et al. 2019).

11.8.3 Hybrid mapping strategies

The preceding architectural subsections treat each architecture in isolation, but real models rarely consist of a single layer type. A vision transformer, for example, combines a patch embedding stage, self-attention layers, and MLP blocks (Dosovitskiy et al. 2021). Those layers create different reuse patterns: the embedding stage can benefit from weight-stationary mapping, attention emphasizes activation movement and tiling, and MLP blocks demand blocked GEMM tiling and fusion. No single dataflow strategy is optimal across all these layers, so hardware mapping becomes hybrid and layer-specific.

Hybrid mapping addresses this heterogeneity by allowing the accelerator to switch strategies at layer boundaries. Each layer presents a different balance of compute intensity, data reuse, and memory access pattern, and the optimal mapping must shift accordingly (Sze et al. 2017). Rather than committing to one dataflow for the entire model, hybrid approaches select weight stationary execution for layers with high weight reuse, activation stationary execution for attention layers with large KV caches, and output stationary execution for layers where minimizing write traffic matters most.

Modern accelerators provide the architectural features needed to realize hybrid mapping in practice. TPU-style systolic arrays, NVIDIA GPUs, and tile-based accelerator designs expose different combinations of local memory, tensor layouts, fusion, and scheduling controls, allowing compilers and runtimes to choose layer-specific strategies rather than one global dataflow for the whole model (Jouppi et al. 2023; NVIDIA Corporation 2020a; Chen et al. 2018). These implementations require programmable memory hierarchies, efficient interconnects, and specialized execution pipelines, reinforcing the hardware-software co-design principle.

However, hybrid mapping remains a design-time optimization. In production workloads, execution conditions change dynamically due to varying input sizes, memory contention, and hardware

³⁸ | **GEMM:** The operation $C = \alpha AB + \beta C$ is the dense linear-algebra primitive that many deep-learning layers lower to. Optimized GEMM libraries such as cuBLAS and oneDNN use register blocking, vectorization, and hierarchical tiling to approach hardware limits on favorable shapes (NVIDIA 2024a; Intel Corporation 2021b). Modern AI accelerators are heavily specialized for GEMM-like tiles: Tensor Cores, systolic arrays, and matrix extensions all exist to accelerate this primitive, which is why GEMM performance is an important predictor of end-to-end throughput across architectures.

resource availability. Machine learning compilers and runtime systems extend these static mapping choices by introducing dynamic scheduling, memory optimizations, and automatic tuning, ensuring that deep learning workloads operate efficiently across diverse accelerators and deployment environments.

11.9 Compiler Support

A single convolution can be implemented with dozens of valid tiling strategies, kernel variants, and memory layouts, most of which perform poorly on a given target. **Machine learning compilers** navigate this complexity by translating dataflow strategies into target-specific executable code. Compiling ResNet-50 for GPU inference exemplifies the process:

1. **Graph optimization** fuses repeated Conv2D-BatchNorm-ReLU patterns into fewer kernels, eliminating intermediate writes that would otherwise consume bandwidth
2. **Kernel selection** chooses Tensor Core implementations for compatible convolutions, exploiting the high arithmetic intensity calculated in the Roofline analysis
3. **Memory planning** determines whether intermediate activations fit in accelerator memory and whether buffers can be reused safely
4. **Computation scheduling** overlaps memory transfers with computation when dependencies allow, hiding part of the transfer latency

In this illustrative scenario, inference time drops from 47 ms to 8 ms, a 5.9× improvement without changing the model’s mathematical function. These numbers are not a ResNet-50 benchmark. They show how graph rewrites, kernel selection, memory planning, and scheduling can translate fusion (section 11.8.1.3) and tiling (section 11.8.1.4) into delivered performance (Chen et al. 2018; NVIDIA 2024c).

This process exemplifies the hardware-software co-design principle established in section 11.1. Machine learning compilers bridge high-level model representations and low-level execution by restructuring computations, selecting kernels, and planning tensor storage (Chen et al. 2018). They build on conventional compiler techniques while adding tensor- and graph-level transformations.

11.9.1 ML compiler design

Machine learning compilers specialize conventional compilation for programs expressed as tensor graphs and operator libraries. General-purpose compilers already optimize loops, vector instructions, parallel code, and memory access; ML compiler stacks add graph rewrites, tensor layouts, shape specialization, accelerator kernel selection, and execution planning (Li et al. 2021).

Table 11.23 contrasts the emphasis of the two layers. The boundary is not absolute: an ML stack often lowers tensor programs into an LLVM, GPU, or vendor backend, so graph-level and instruction-level compilation cooperate.

Table 11.23: Compiler Optimization Priorities: ML compiler stacks add tensor- and graph-level decisions to conventional instruction- and loop-level compilation. In a production toolchain, the two layers commonly cooperate rather than replace one another.

Aspect	Traditional compiler	Machine learning compiler
Input representation	Source or intermediate-representation program	Tensor graph or tensor-level intermediate representation
Execution Model	Control flow, loops, threads, vectors, and tasks	Tensor operators, kernels, graphs, and accelerator streams
Optimization priorities	Instruction selection, loop transforms, vectorization, registers	Graph rewrites, fusion, layouts, shapes, and kernel selection
Memory management	Objects, stack/heap, caches, locality, and prefetching	Tensor lifetimes, buffer reuse, layouts, tiling, and device transfers
Target Hardware	CPUs, GPUs, and other programmable targets	CPUs, GPUs, TPUs, and custom accelerators
Compilation output	Machine code or lower-level intermediate representation	Kernels plus a hardware-specific graph or execution plan

The table explains why compiler configuration can change performance even when model code is unchanged. ML compilers own the hidden layer that maps graph-level tensor operations onto hardware-specific kernels, layouts, and schedules; when that mapping is poor, the model leaves arithmetic units idle or moves the same bytes repeatedly.

Systems Perspective 11.2: The hidden optimization layer

Most practitioners never interact directly with ML compilers, yet compiler quality often determines whether a model achieves a low or high fraction of hardware peak performance. Calling `model.compile()` in Keras, `torch.compile()` in PyTorch, or deploying through TensorRT invokes multi-stage optimization pipelines. These pipelines perform four hidden optimizations:

- **Fuse operations** never explicitly combined by the developer (Conv2D + BatchNorm + ReLU → single kernel)
- **Reorder computations** to improve memory locality (tiling large matrix multiplies)
- **Select kernels** from libraries containing hundreds of hand-tuned implementations
- **Transform tensor layouts** between what the model definition expects and what hardware prefers

This matters practically: the same model definition can run substantially faster or slower depending on compilation backend, graph lowering, kernel selection, and runtime configuration. When performance does not meet expectations, compiler configuration and backend selection are often the first optimization levers, requiring no changes to model architecture or training procedure.

11.9.2 ML compilation pipeline

Machine learning models, as defined in modern frameworks, are initially represented in a high-level computation graph that describes operations on tensors. These representations must be lowered into executable code for target hardware such as CPUs, GPUs, TPUs, and custom AI chips. An ML compilation pipeline performs this transformation (Chen et al. 2018; Google 2025; Lattner et al. 2020).

A useful compilation model has five responsibilities, although real systems interleave them. Graph optimization restructures the computation. Kernel selection maps operations to target implementations. Memory planning assigns layouts, buffers, and lifetimes. Scheduling orders kernels and transfers subject to dependencies. Code generation then emits or packages the target-specific executable components.

At each stage, the compiler applies the optimizations developed in section 11.8: kernel fusion, tiling, data movement strategies, and computation placement. These optimizations are systematically incorporated into the final execution plan, which is why machine learning acceleration depends as much on compiler-driven software optimization as on hardware improvements.

11.9.3 Graph optimization

AI accelerators provide specialized hardware to speed up computation, but raw model representations are not inherently optimized for execution on these accelerators. Machine learning frameworks define models using high-level computation graphs, where nodes represent operations (such as convolutions, matrix multiplications, and activations), and edges define data dependencies. However, if executed as defined, these graphs often contain redundant operations, inefficient memory access patterns, and suboptimal execution sequences that can prevent the hardware from operating at peak efficiency.

For example, transformer self-attention can create large intermediate score and probability matrices. A naïve implementation that materializes and rereads those intermediates from high-bandwidth memory pays excessive memory traffic, while IO-aware attention kernels tile the computation through fast memory to avoid that traffic (T. Dao et al. 2022). Similarly, in a CNN, applying batch normalization and activation functions as separate operations after each convolution leads to unnecessary intermediate memory writes, increasing memory bandwidth usage. These inefficiencies are

addressed during graph optimization, where the compiler restructures the computation graph to eliminate unnecessary operations and improve memory locality (Chen et al. 2018).

Graph optimization transforms this high-level computation graph into an optimized execution plan before hardware mapping. Rather than requiring manual optimization, the compiler systematically applies transformations that improve data movement, reduce redundant computations, and restructure operations for efficient parallel execution (Chen et al. 2018; Jia et al. 2019). At this stage, the compiler works at a hardware-agnostic level, focusing on high-level restructuring before hardware-specific optimizations are applied in subsequent compilation phases (section 11.9.6.2).

Graph optimization first removes traffic that the high-level graph representation would otherwise create. Kernel fusion merges consecutive operations to eliminate unnecessary memory writes and reduce the number of kernel launches, which is particularly effective in convolutional neural networks where convolution, batch normalization, and activation functions appear in fixed sequences. Computation reordering adjusts execution order to improve data locality and parallel execution; in transformer models, this reordering enables reuse of cached key-value pairs rather than repeated memory reloads.

Redundant computation elimination serves the same goal from the compute side. By identifying and removing duplicate or unnecessary operations, the compiler avoids repeated work in models with residual connections where common subexpressions might otherwise be recomputed. Memory-aware dataflow adjustments then refine tensor layouts and optimize movement; for example, tiling matrix multiplications to meet the structural requirements of systolic arrays in TPUs aligns the graph with the accelerator's strengths. Together, these techniques prepare the model for acceleration by minimizing overhead and balancing computation against memory resources.

Modern AI compilers implement these rewrites through automated pattern recognition and structured rules, but the core responsibilities are the same across compiler stacks: find fusible patterns, choose layouts that match the target memory hierarchy, and preserve the model's mathematical meaning while exposing hardware-specific optimization opportunities. XLA, TVM, TensorRT, and MLIR are representative systems that emphasize different target constraints, from graph-level fusion to tensor-layout search and multi-stage lowering. The systems lesson is not the product list; it is that compiler restructuring turns a framework graph into an execution plan the accelerator can sustain. Without this restructuring, a large transformer model on an edge device may suffer excessive memory stalls; with it, reduced bandwidth consumption and latency can make real-time inference feasible on resource-constrained devices.

After graph-level rewrites expose the available operations, kernel selection determines which target-specific implementation should realize each one. Later memory and scheduling decisions may feed back into that choice.

11.9.4 Kernel selection

Kernel selection turns the optimized graph into a hardware contract. A kernel is a specialized implementation of a computational operation designed to run efficiently on a particular hardware architecture. Most accelerators, including GPUs, TPUs, and custom AI chips, provide multiple kernel implementations for the same operation, each optimized for different execution scenarios. Choosing the right kernel determines whether the accelerator maximizes computational throughput, avoids memory stalls, and keeps specialized processing elements busy (Chen et al. 2018; Zheng et al. 2020).

Kernel selection builds upon graph optimization, mapping the structured execution plan to the most efficient implementation available for each operation. Poor kernel choices can nullify the benefits of prior optimizations by introducing unnecessary computation overhead or memory bottlenecks (Chen et al. 2018).

Transformer matrix multiplications illustrate the choice. A CPU backend may select a vectorized library kernel, a GPU backend may use Tensor Cores when shape and precision permit, and a TPU backend maps compatible tiles to its matrix unit. An INT8 inference kernel is useful only when the model has been quantized and validated at that precision; lower precision changes both performance and numerical behavior.

In many cases, the compiler's decision is which existing implementation to trust rather than whether to generate a kernel from scratch. cuDNN and cuBLAS offer optimized kernels for deep

learning on NVIDIA GPUs, oneDNN provides optimized execution for Intel architectures, ACL (Arm Compute Library) targets Arm-based devices, and Eigen and BLIS provide efficient CPU-based implementations. These libraries encode hardware-specific knowledge so the compiler can choose a preoptimized kernel rather than reinventing an execution strategy for each platform.

AI compilers use heuristics,³⁹ profiling, and cost models to decide among these options. The selection method depends on how much uncertainty the compiler can tolerate before execution begins.

Rule-based selection applies predefined heuristics based on known hardware capabilities. For instance, XLA, the compiler used in TensorFlow, automatically selects Tensor Core-optimized kernels for NVIDIA GPUs when mixed-precision execution is enabled. These predefined rules allow fast, reliable decisions without extensive analysis.

Profile-guided selection pays more search cost to reduce uncertainty. TVM uses AutoTVM to benchmark kernel options empirically and tune execution strategies based on real execution times, so operations are assigned to implementations that perform well under actual deployment conditions.

Cost model-based selection estimates execution time and memory consumption before profiling every option. MLIR provides compiler infrastructure in which target-specific passes can implement this kind of selection (Lattner et al. 2020). A concrete compiler can model how candidate kernels interact with the accelerator's compute units and memory hierarchy, then select a kernel that minimizes estimated execution cost.

Precision-aware selection adds the numerical constraint to the same decision. Training workloads often prioritize FP32 or BF16 to maintain model accuracy, whereas inference workloads favor FP16 or INT8 to increase speed and reduce power consumption. For example, an NVIDIA GPU running inference with TensorRT can select among calibrated FP16 and INT8 engine profiles that were built for the model's accuracy constraints and input shapes. This trade-off between precision and performance is a key aspect of kernel selection, especially in resource-constrained environments.

Some compilers extend selection into adaptive tuning, where execution strategies adjust to workload and resource conditions. AutoTVM in TVM measures kernel performance across workloads and refines execution strategies; TensorRT applies optimized engine profiles based on batch size, memory constraints, and supported precision; Google's TPU compiler specializes execution plans for the target TPU topology and shape profile. The consequences of poor kernel selection are significant: a transformer model assigned a nontensor-core kernel for matrix multiplications may execute at only a fraction of possible performance, while a model designed for FP32 execution may lose accuracy if forced onto an INT8-optimized kernel. Kernel selection is therefore as much about numerical correctness as performance.

After an initial kernel selection, memory planning and scheduling determine buffer lifetimes, workspaces, launch order, and legal overlap. These decisions can also cause the compiler to revisit a kernel whose workspace or layout no longer fits the emerging plan.

11.9.5 Memory planning

The memory planning phase ensures that data is allocated and accessed in a way that minimizes memory bandwidth consumption, reduces latency, and maximizes cache efficiency (Roesch et al. 2018; Chen et al. 2018). Even with the most optimized execution plan, a model can still suffer from severe performance degradation if memory is not managed efficiently.

Machine learning workloads are memory-intensive, requiring frequent movement of large tensors between different levels of the memory hierarchy. The compiler must determine how tensors are stored, how they are accessed, and how intermediate results are handled to prevent memory from becoming the bottleneck.

The memory planning phase optimizes tensor layouts, memory access patterns, and buffer reuse to prevent unnecessary stalls and memory contention during execution. Tensors are arranged in formats that align with hardware access patterns, minimizing format conversions. Memory accesses are structured to reduce cache misses and stalls, lowering overall bandwidth consumption. Buffer reuse reduces redundant memory allocations by managing intermediate results so that completed buffers are reclaimed promptly. Together, these strategies ensure that data is efficiently placed and accessed, enhancing both computational performance and energy efficiency.

³⁹ | **Heuristic in kernel selection:** A practical rule that chooses a promising implementation without exhaustively benchmarking the full candidate space. Tile sizes, layouts, precision modes, fusion choices, and shape constraints can create many legal GEMM variants. Heuristics reduce tuning cost but may miss a faster candidate, which is why autotuners such as AutoTVM profile selected options on the target hardware.

Balancing memory availability, reuse, and access efficiency across multiple hierarchy levels makes memory planning one of the most complex compiler problems. AI compilers use several strategies to manage memory effectively and prevent unnecessary data movement.

Tensor layout optimization determines how tensors should be arranged in memory to maximize locality and prevent unnecessary format conversions. As section 11.8.1.2 established, different hardware accelerators favor different physical layouts depending on the backend kernel and precision mode. NVIDIA's cuDNN convolution path historically expected NCHW for many FP32 kernels, while the channels-last NHWC layout aligns with Tensor Core memory coalescing for FP16 and INT8 paths; TensorFlow/XLA may choose internal layouts during lowering for the target backend. Compiler and library stacks transform tensor layouts based on the kernel and precision selected for the target hardware, ensuring that memory accesses are aligned for maximum efficiency (NVIDIA Corporation 2021b; Google 2025).

Buffer allocation and reuse complements layout optimization: the compiler minimizes memory footprint by reusing intermediate storage whenever possible. Deep learning workloads generate many temporary tensors, such as activations and gradients, which can quickly overwhelm on-chip memory if not carefully managed. Instead of allocating new memory for each tensor, the compiler analyzes the computation graph to identify opportunities for buffer reuse, ensuring that intermediate values are stored and overwritten efficiently (Roesch et al. 2018).

Minimizing data movement between hierarchy levels is equally critical. AI accelerators typically have a mix of high-speed on-chip memory (such as caches or shared SRAM) and larger, but slower, external DRAM. If tensor data is repeatedly moved between these memory levels, the model may become memory bound, reducing computational efficiency. To prevent this, compilers use tiling strategies that break large computations into smaller, memory-friendly chunks, allowing execution to fit within fast, local memory and reducing the need for costly off-chip memory accesses. The consequences of neglecting memory planning are concrete: a CNN running on a GPU may achieve high computational efficiency in theory, but if its convolutional feature maps are stored in an incompatible layout that necessitates repeated format conversions, the resulting overhead can negate the gains from graph optimization and kernel selection entirely. With memory allocation determined, the compiler must next decide when and where each computation executes.

11.9.6 Computation scheduling

With graph optimization completed, kernels selected, and memory planning finalized, computation scheduling determines the execution order and resource assignment for each operation. This phase determines when and where each computation should be executed, ensuring that workloads are efficiently distributed across available processing elements while avoiding unnecessary stalls and resource contention (Zheng et al. 2020).

Without effective scheduling, massive parallelism goes to waste: computational units sit idle, memory bandwidth goes underutilized, and execution efficiency degrades. Computation scheduling keeps processing elements active, manages execution dependencies correctly, and distributes workloads across the hardware schedule space (Chen et al. 2018; Zheng et al. 2020).

The scheduling phase coordinates parallel execution, synchronization, and resource allocation. Task partitioning decomposes computations into units that can be distributed among multiple compute cores. Execution order optimization determines the sequence for launching operations, maximizing hardware performance while reducing stalls. Resource allocation and synchronization ensure that compute cores, memory bandwidth, and shared caches are used without contention.

11.9.6.1 Implementation in AI compilers

Scheduling strategies are highly dependent on the underlying hardware architecture, since different AI accelerators have unique execution models. AI compilers implement several strategies to optimize scheduling for efficient execution.

Task partitioning divides large computational graphs into smaller units that can execute in parallel. On GPUs, this typically means mapping matrix multiplications and convolutions to thousands of CUDA cores, while on TPUs, tasks are partitioned to fit within systolic arrays that operate on structured data flows (Norrie et al. 2021). In CPUs, partitioning is often focused on breaking computations into vectorized chunks that align with SIMD execution. In each case, the goal is to keep every core active throughout execution.

Beyond task partitioning, scheduling involves optimizing execution order to minimize dependencies and maximize throughput. Many AI models include operations that can be computed independently (for example, different batches in a batch processing pipeline) alongside operations that have strict dependencies (for example, recurrent layers in an RNN). AI compilers analyze these dependencies and attempt to rearrange execution where possible, reducing idle time and improving parallel efficiency. In transformer attention, IO-aware kernels make this scheduling problem concrete by loading blocks of queries, keys, and values into fast memory, using them while resident, and evicting them in an order that reduces high-bandwidth-memory traffic (T. Dao et al. 2022).

Resource allocation and synchronization determine how compute cores share memory and coordinate execution. Modern AI accelerators often support overlapping computation and data transfers, meaning that while one task executes, the next task can begin fetching its required data. Compilers take advantage of this by scheduling tasks in a way that hides memory latency, ensuring that execution remains compute bound rather than memory-bound (Chen et al. 2018). In production inference stacks, optimized runtimes and compiler-generated schedules coordinate kernel launch order, stream execution, and synchronization so the accelerator does not stall between dependent kernels (NVIDIA 2024c; Zheng et al. 2020). Poor scheduling decisions can negate the benefits of all prior compilation phases: a CNN with highly optimized kernels and efficient memory layouts will still suffer reduced throughput if compute units remain idle between kernel launches, and a transformer on a TPU may underperform if attention layers are not scheduled to overlap with memory transfers.

11.9.6.2 Code generation

With scheduling complete, the final compilation stage translates this optimized execution plan into hardware-specific instructions. Unlike the previous phases, which required AI-specific optimizations, code generation follows many of the same principles as traditional compilers. This process includes instruction selection, register allocation, and final optimization passes, ensuring that execution makes full use of hardware-specific features such as vectorized execution, memory prefetching, and instruction reordering. Crucially, however, instruction selection for ML targets is not generic: the compiler must emit instructions that engage the hardware’s matrix-specific ISA extensions. On NVIDIA GPUs, this means emitting Parallel Thread Execution (PTX) instructions such as `mma.sync.aligned` to invoke Tensor Cores directly, as shown in listing 11.14. On Intel CPUs with Advanced Matrix Extensions (AMX), the compiler targets tile-multiply instructions operating on 2D register tiles. On Arm CPUs with the Scalable Matrix Extension, the target is outer-product accumulation across scalable matrix tiles. A code generation backend that emits generic floating-point instructions instead of these extensions leaves the hardware’s primary matrix engines idle, which can reduce effective throughput by an order of magnitude regardless of how well prior compilation phases performed. For CPUs and GPUs, AI compilers typically generate machine code or optimized assembly instructions, while for TPUs, field-programmable gate arrays (FPGAs),⁴⁰ and other accelerators, the output may be optimized bytecode or execution graphs that are interpreted by the hardware’s runtime system.

11.9.7 From compilation to runtime

The compiler transforms a high-level model into an execution plan tailored to target hardware, but that plan still embeds assumptions about future conditions: shapes, workspace, device capabilities, and sometimes concurrency. Graph optimization, kernel selection, memory planning, and scheduling can make execution efficient; they cannot guarantee full utilization under every runtime workload.

Ahead-of-time compiler and engine-building optimizations occur before execution begins. This static nature enables aggressive graph optimization but limits adaptation when input shapes or resource conditions diverge from the prepared profiles. Graph restructuring, fusion, tactic selection, tiling, precision choices, and much of memory planning are therefore based on the shapes and hardware known at build time. Just-in-time compilation systems may compile additional variants later, but each generated kernel still has a fixed implementation.

Production AI systems inhabit a dynamic world that may not match these static assumptions. Batch sizes vary, multiple workloads may compete for an accelerator, and thermal throttling can reduce

40 | **Field-programmable gate array:** “Field-programmable” means the logic fabric is configurable after manufacturing, contrasting with fixed-function ASICs. FPGAs can improve performance for latency-sensitive data center services by implementing custom pipelines matched to a particular workload (Putnam et al. 2014). This reconfigurability makes FPGAs attractive for rapidly evolving ML architectures where committing to an ASIC risks obsolescence, but the requirement for hardware description languages (Verilog/VHDL) and compilation times measured in hours creates a productivity barrier that limits adoption to deployments where the efficiency benefit justifies the engineering cost.

sustained performance below short-benchmark peaks. Runtimes and serving systems respond by selecting among supported profiles, managing buffers and streams, batching requests, and applying admission control; they do not generally invent new fusion or tiling strategies for a built engine. The serving chapter treats batching, admission control, and service-level objectives as end-to-end system problems (chapter 13); here the narrower question is how the accelerator runtime dispatches prepared execution plans once a request or batch reaches the hardware.

11.10 Runtime Support

AI runtimes execute compiled graphs and engines while managing input-dependent shapes, buffers, streams, and supported execution profiles; TensorRT is one representative production inference stack. TensorRT’s builder compiles the network and selects tactics when creating an engine, while its runtime supplies inputs, selects a compatible optimization profile, and dispatches the prepared implementation (NVIDIA Corporation 2026a, 2026b).

AI runtimes manage three interrelated aspects of execution. First, they dispatch kernels or prepared graph variants that are legal for the current shapes and device. Second, they bind and reuse tensor buffers, workspaces, streams, and transfer resources. Third, higher-level framework runtimes can participate in multi-device execution, as in pipeline-parallel systems such as GPipe (Huang et al. 2019) and device-placement methods (Mirhoseini et al. 2017). The low-level inference runtime, training framework, and serving scheduler may divide these responsibilities differently.

AI runtimes complement compiler-based optimizations by handling these execution aspects dynamically. Comparing AI runtimes to traditional software runtimes clarifies why machine learning workloads require specialized execution strategies.

11.10.1 ML runtime architecture

General-purpose runtimes manage threads, tasks, objects, memory, and asynchronous I/O across diverse programs. AI runtimes build on those mechanisms but represent work in terms of tensors, kernels, graphs, device buffers, and accelerator queues.

Table 11.24 contrasts their emphasis. AI runtimes must respect graph dependencies and tensor shapes while coordinating kernels and transfers, but they coexist with conventional language, operating-system, and device runtimes.

Table 11.24: Runtime Execution Models: AI runtimes add tensor-, graph-, and accelerator-aware resource management to conventional runtime services. The exact boundary among framework, compiler runtime, device library, and serving scheduler varies by stack.

Aspect	Traditional runtime	AI runtime
Execution Model	Functions, threads, tasks, events, and asynchronous I/O	Tensor graphs, kernels, streams, and device events
Task scheduling	Threads and tasks across processor resources	Dependency-aware kernel and transfer dispatch
Memory Management	Objects, stacks, heaps, pools, and virtual memory	Tensor buffers, workspaces, device memory, and buffer reuse
Optimization Priorities	Responsiveness, throughput, locality, and resource sharing	Shape compatibility, transfer overlap, reuse, and utilization
Adaptability	Dynamic scheduling and allocation within program semantics	Runtime choices within compiled graphs and available kernels
Target Hardware	CPUs and heterogeneous devices	CPUs, GPUs, TPUs, and custom accelerators

Tensor memory is a major specialization. AI runtimes often combine static lifetime plans for fixed shapes with pools or dynamic allocation for variable shapes, then bind those buffers to device kernels and transfers. Poor planning can increase peak memory, prevent buffer reuse, or add transfers; cache behavior remains largely under the compiled kernels and hardware.

AI runtimes support variability within the bounds of their compiled engines or available kernel libraries. They can manage dynamic tensor buffers, select a compatible shape profile, coordinate streams, and participate in multi-device execution. Request batching and fleet-level resource allocation usually belong to the serving scheduler rather than the low-level accelerator runtime.

Production conditions can differ from the benchmark environment. The A100 SXM has a maximum thermal design power (TDP) of 400 W, while the A100 PCIe has a 300 W maximum TDP; these are different form factors with different cooling envelopes, so a deployment may not sustain the throughput measured on another variant. Batch size, memory availability, and resource contention can also vary. The system must therefore prepare appropriate engine profiles and use serving controls such as batching and admission control rather than assume that the runtime will retune compiled kernels in response to temperature or load.

To see how these runtime mechanisms work together, consider a transformer inference request arriving at a production server. The serving system forms a supported batch, the runtime selects a compatible prebuilt profile and manages its buffers, and the engine dispatches the kernels chosen during compilation or engine construction. The subsections that follow examine these bounded runtime choices using the request as a running example.

11.10.2 Dynamic kernel execution

While static compilation provides a foundation, efficient execution may require choosing among prepared variants. When a transformer request arrives, its batch and sequence dimensions determine which supported optimization profile and kernels can execute it. Shapes outside those profiles require a different engine or additional compilation rather than arbitrary runtime retuning.

Individual operations are assigned implementations during compilation or engine construction. At execution time, the runtime binds buffers, selects legal variants when the engine provides them, and launches work in dependency order.

The same constraint appears in image classification. If an incoming batch of high-resolution images falls outside the memory assumptions of a prepared profile, the serving system must select a compatible profile or engine, split the batch, or rebuild the engine. The runtime cannot invent a new tiling strategy for an already built engine.

For the running transformer inference request, sequence length may vary between calls. A plan optimized for one fixed length can underutilize compute resources on shorter sequences or exceed its workspace assumptions on longer ones. A runtime can select a compatible prepared profile and allocate its required buffers; unsupported lengths require another compiled variant or a more general kernel.

Overlapping computation with memory movement can mitigate transfer bottlenecks. When dependencies, page-locked buffers, streams, and copy engines permit concurrency, asynchronous execution and double buffering transfer batch $n+1$ while computing batch n . This can hide part of the host-to-device transfer latency, but it does not eliminate stalls when transfer time exceeds compute time or resources serialize.

Convolutional layers illustrate this boundary: TensorRT may fuse operations and benchmark convolution tactics while building the engine. During inference, the runtime dispatches those prepared layers; the GPU's hardware schedulers place their thread blocks on available SMs.

Selecting among prepared execution strategies in response to request shape and system conditions can improve both training and inference performance. This adaptation depends on having the right kernel available. In the transformer inference scenario, the runtime must choose among implementations prepared for the request's supported shape and precision.

11.10.3 Runtime kernel selection

While compilers perform an initial selection of kernels based on static analysis, AI runtimes may still choose among precompiled or library-provided variants during execution. Real-time factors, such as available memory, hardware utilization, and workload priorities, may differ from the assumptions made during compilation. In the transformer scenario, the compiler and framework determine the legal precision paths, while the runtime selects the kernel variant that best fits the current sequence length, batch shape, and available hardware resources. Runtime selection adapts execution to changing conditions, but it remains bounded by the numerical formats and kernels the model has been prepared to use.

For instance, transformer-based language models spend substantial time in matrix multiplications. Mixed-precision systems such as Megatron-LM use FP16 execution on GPU Tensor Cores to increase

throughput (Shoeybi et al. 2019). Which operations remain in FP32 is a model-conversion or engine-building decision validated before deployment; the runtime does not infer numerical instability and change precision on its own.

Memory constraints also influence kernel selection. When memory bandwidth is limited, the runtime may select a prepared kernel or profile whose tiling and workspace requirements fit the available resources. Creating a new tiling strategy requires engine building or compilation rather than an arbitrary runtime adjustment.

Batch size also influences kernel selection. When a stack provides multiple compatible variants, the runtime may use a latency-oriented implementation for small batches and a throughput-oriented one for larger batches. The benefit depends on the prepared kernel set and the cost of switching profiles. The final pipeline stage then determines when each selected kernel runs.

11.10.4 Kernel scheduling and utilization

Kernel dispatch completes the runtime pipeline. Returning to the transformer request, the runtime launches the prepared attention, normalization, and activation kernels on streams while respecting graph dependencies. On a GPU, hardware block schedulers assign thread blocks to SMs; the runtime does not directly distribute individual operations across cores. Other accelerators expose analogous queues or execution commands (Jouppi et al. 2017).

In image recognition models, independent kernels may overlap on separate streams when their dependencies and resource use permit. Within each GPU kernel, however, the compiled grid defines the work and GPU hardware schedules its thread blocks; the runtime does not distribute individual filters across processing units.

Memory management reinforces the same goal through buffer reuse, asynchronous transfers, and explicit staging where the API supports it. Hardware caches remain hardware-managed; a runtime cannot generally place arbitrary intermediate tensors into cache.

Together, profile selection, buffer management, and kernel dispatch form the runtime pipeline. For the transformer request, the runtime chooses among implementations and shapes prepared by the compiler or engine builder, then launches them efficiently. Changes to fusion, tiling, or legal precision generally require rebuilding or recompiling rather than continuous runtime retuning.

The compiler and runtime systems examined thus far optimize execution within single accelerators, but the largest AI workloads exceed what any single chip can deliver. Single-chip optimizations can achieve impressive results through compiler optimization, dataflow selection, fusion, and memory planning. Yet for the largest AI workloads, even well-optimized single-chip execution proves insufficient.

Consider the estimated 3.14×10^{23} floating-point operations associated with training GPT-3 (Brown et al. 2020). As a scale comparison, divide that historical operation count by a later H100's peak FP8 rate of 1.98 PFLOP/s. The arithmetic lower bound is about 5 years on one device, and 8.4–12.6 years under the illustrative 40–60 percent utilization range (Choquette 2023). This is not a reconstruction of GPT-3 training; the precision, hardware, communication, and workload are counterfactual. It simply shows why model-scale compute or service-scale request volume can exceed the useful capacity of one accelerator.

11.11 Multi-Chip Scaling

A single H100 has a large FP8 peak, yet a training-time or throughput target may still require many accelerators. The techniques covered earlier remain the foundation for multi-chip scaling; each device still needs an efficient dataflow. The new lesson is that communication adds further hierarchy boundaries. Moving data through on-chip storage, HBM, an intra-node link, and a cluster fabric generally increases latency and energy, so scaling is not merely a question of adding chips.

When single-accelerator capacity proves insufficient, the design problem shifts from feeding one chip to choosing which communication boundary the workload can tolerate. Practitioners encounter these boundaries in production even when most optimization work remains inside a single accelerator.

11.11.1 Multi-chip scaling approaches

Large AI systems scale beyond individual accelerators by moving the communication boundary outward, and each boundary changes the dominant trade-off. The sequence begins inside the package. Chiplet-based architectures partition large designs into smaller, modular dies interconnected within one package, bypassing manufacturing limits of monolithic chips while preserving relatively low communication latency. The next boundary is the node: multi-accelerator servers connect several chips through board- or server-level interconnects. Each accelerator has dedicated memory and compute resources, so workloads split through data parallelism (each accelerator processes different batches) or model parallelism (different accelerators handle different network layers). High-bandwidth intra-node interconnects can enable efficient gradient synchronization, though realized performance depends on topology and collective communication efficiency.

Beyond the node, the boundary expands into the cluster. Purpose-built data center fabrics coordinate hundreds of accelerators, making topology and collective communication algorithms central determinants of scaling efficiency; near-linear scaling is achievable on some workloads when communication overhead is controlled. Wafer-scale integration is the counter-move: instead of pushing the boundary outward, it collapses more computation back into one large device. Platforms such as Cerebras WSE-class systems integrate extremely large numbers of transistors and cores on a single device, reducing inter-chip communication overhead while introducing their own challenges in thermal dissipation, fault tolerance, and manufacturing yield.

11.11.2 Why scaling introduces new constraints

The transition from single-chip to multi-chip architectures introduces communication overhead and other constraints that reshape optimization. Communication, synchronization, load balance, and parallel work per device can all limit scaling. Amdahl's Law⁴¹ provides one first-order model for exposed work that does not shrink with additional devices. For hundred-billion-parameter models, an unsharded full-gradient payload can occupy hundreds of gigabytes per step, and an AllReduce⁴² must aggregate it; precision, sharding, compression, and the collective algorithm change the bytes each device transfers.

The first-order quantity is the gradient payload. For a model with P parameters, that payload is roughly the parameter count multiplied by the bytes stored for each gradient element before optimizer state, padding, and protocol overhead enter. Scaling therefore improves only when the saved compute time exceeds the time to move this payload through the chosen interconnect and collective algorithm.

This overhead explains why large accelerator counts can show diminishing returns unless the system reduces exchanged data, overlaps communication with useful computation, or chooses a better parallelization pattern. The memory model also changes: separate accelerator memories are commonly managed explicitly, and collective operations establish the required synchronization rather than presenting one fully coherent memory image across the cluster. Coherent links exist in some systems, but their scope and cost are architecture-dependent.

Once computation spans many links, chips, and memory stacks, reliability and energy become part of the same scaling story. Large-scale systems must handle component failures gracefully because the probability of at least one failure rises with system size. For this chapter, the hardware lesson is enough: distributed TPU systems must tolerate component failures at system scale (Jouppi et al. 2023), while Cerebras uses redundant cores and fabric links to replace manufacturing-defective cores and restore the wafer's logical mesh (Lie 2021). Data movement also grows more expensive with distance, transforming distributed training into a careful balance between computation parallelism and communication efficiency.

Data center scaling and edge deployment represent opposite ends of a deployment spectrum, yet they share the same core principles. Data center scaling coordinates many high-throughput accelerators, while edge scaling fits useful AI into a few constrained watts. Both cases require matching workload characteristics to hardware capabilities while minimizing data movement. The principles of compute specialization, memory hierarchy optimization, and workload mapping apply at both scales; only the constraints differ. Data centers optimize for aggregate throughput within power budgets measured in megawatts; edge devices optimize for responsiveness within tight

⁴¹ | **Amdahl's law (scaling limit):** section D.2.3 formalizes the bound for a fixed nonaccelerated fraction. If 5 percent of a fixed workload remains exposed and unaffected by device count, that model caps speedup at $20\times$. Distributed training is more complicated because communication can overlap computation and its cost can itself change with scale, but the example explains why exposed synchronization must be reduced or hidden.

⁴² | **AllReduce:** A collective operation from MPI that aggregates values across processes (the "reduce") and distributes the result back to every process (the "all"). In this chapter, the term appears only to identify why accelerator-to-accelerator bandwidth matters for training workloads. At scale, algorithms, topology choices, and runtime protocols determine how costly this synchronization becomes.

battery and thermal envelopes. The same vision model that runs comfortably in a data center may need a radically different mapping strategy on a smartphone or microcontroller.

11.12 Heterogeneous SoC Design

Mobile, automotive, and IoT deployments often face tighter per-device power, thermal, and latency limits than data center hardware. Heterogeneous SoCs integrate CPU cores, GPU shaders, digital signal processors (DSPs), and neural processing units (NPUs) under shared memory and system budgets. The mapping problem remains familiar, but it now includes operator support, transfer overhead, real-time deadlines, and thermal state.

11.12.1 Mobile SoC architecture evolution

Modern mobile AI engines exemplify heterogeneous computing by coordinating CPU cores, GPU shaders, DSPs, and dedicated NPUs⁴³ across a shared memory hierarchy. Workload distribution lets computer vision kernels execute on GPU or NPU paths, audio processing use DSP arithmetic units, and matrix-heavy neural-network layers use NPU-optimized engines when the operator set is supported. This coordination requires careful scheduling to meet real-time constraints while managing thermal throttling and battery life.



Example 11.5: Heterogeneous microcontrollers

Scenario: An embedded smart doorbell must run periodic person detection within a tight memory, energy, and latency budget.

Diagnosis: On a supported deployment, a general-purpose Cortex-M path misses the measured energy or latency target. A micro-NPU such as Arm Ethos-U can offload compatible convolution operators, while the CPU retains control flow and unsupported operations.

Systems lesson: Micro-NPU specialization can improve the energy and latency of supported operators, but feasibility must be measured for the whole pipeline, including fallbacks and sensor processing.

⁴³ | **NPU (neural processing unit):** The NPU's specialized matrix engines are optimized for dense tensor operations, providing the hardware basis for the workload distribution described. This specialization creates a critical constraint for the scheduler: any AI operator not mapped to the NPU's supported data paths must "fall back" to a GPU or CPU. This fallback can erase the NPU's energy-efficiency advantage, complicate real-time latency budgets, and contribute to thermal pressure on mobile devices.

Some mobile SoC designs emphasize diverse processor specialization, while vertically integrated strategies highlight how tight hardware-software co-design can enable tightly coordinated heterogeneous execution. Unified memory architectures can reduce explicit data copying overhead, and different compute blocks can be scheduled for different operator types (for example, matrix-heavy layers on an NPU, convolutional operators on a GPU, and control flow on the CPU). This coordination supports interactive on-device experiences, though realized latency depends on the full pipeline and device thermal conditions.

Beyond vertically integrated solutions, IP licensing models allow SoC designers to customize processor combinations based on target applications, mixing CPU, GPU, DSP, and NPU blocks. This modular flexibility allows automotive SoCs to emphasize deterministic real-time processing while smartphone SoCs optimize for interactive performance and battery efficiency.

11.12.2 Strategies for dynamic workload distribution

With multiple specialized processors available on heterogeneous SoCs, the critical challenge becomes intelligently distributing neural network operations across these resources to maximize performance while respecting power and latency constraints. Consider a concrete example: an engineer deploying a real-time object detection pipeline on a mobile SoC with a CPU, GPU, and NPU. The pipeline has three stages: a MobileNet backbone for feature extraction, nonmaximum suppression (NMS) for postprocessing, and a display overlay for rendering bounding boxes. The backbone consists of depthwise separable convolutions with regular, predictable access patterns and low compute cost, making it a good fit for an NPU when the operator set is supported, even though depthwise layers are often memory-bound rather than high-arithmetic-intensity kernels. NMS, by contrast, involves conditional branching over variable-length candidate lists, with irregular memory access that maps poorly to the NPU's fixed dataflow. The CPU handles NMS more efficiently because its branch predictor and large caches accommodate the unpredictable control flow. Finally, the

display overlay involves pixel-level compositing across the entire frame, a massively parallel but arithmetically simple workload that maps naturally to the GPU's shader cores. This three-way split, NPU for the backbone, CPU for NMS, GPU for the overlay, achieves lower latency and lower power than running the entire pipeline on any single processor.

The example shows how supported operators, measured transfer cost, and the current system budget determine pipeline partitioning. Regular convolutions often map well to GPU or NPU paths, while irregular control flow may favor a CPU. Attention performance depends on shape, precision, kernel availability, and setup overhead; sequence length alone does not determine the best processor.

The best supported assignment can also change with operating mode. A system may reduce frame rate, select a smaller model, lower accelerator frequency, or choose another validated backend as battery, thermal state, and concurrent load change. Moving work is not inherently more efficient because transfers and fallback kernels can outweigh the gain. Safety-critical applications further prioritize bounded latency and validated behavior over peak throughput.

Compounding the processor selection challenge, shared memory architectures require arbitration when multiple processors access LPDDR simultaneously. Mobile memory controllers may prioritize real-time camera or display paths over background AI tasks, forcing neural-network runtimes to adapt their execution patterns to available bandwidth. This arbitration becomes critical during memory-intensive operations like large language model inference, where parameter streaming from DRAM must be carefully coordinated across processors.

11.12.3 Power and thermal management

Mobile AI workloads must maintain high performance while operating within strict power budgets and thermal envelopes. These constraints require tight coordination across heterogeneous processors.

Heterogeneous SoCs use **dynamic voltage and frequency scaling (DVFS)** across power domains to manage performance within a shared envelope. Raising one domain's operating point can leave less thermal or electrical headroom for others. Control policies react to utilization, deadlines, temperature, and platform-specific power limits; transitions themselves have latency and stability costs.

When DVFS alone cannot maintain the power envelope, mobile SoCs implement thermal throttling through a mixture of frequency reduction, model adaptation, and task migration. When the NPU approaches thermal limits during intensive neural network processing, the runtime can shift selected operators to another supported processor, lower inference frequency, or choose a smaller model profile. This approach preserves service availability during thermal events, though it requires detailed workload characterization to predict execution time and power consumption across different processors.

Mobile AI systems may also adapt quality, model choice, or inference frequency to battery and charging state. The most energy-efficient processor is workload- and implementation-dependent, so a lower-power operating mode must be selected from measured, validated alternatives rather than inferred from processor labels.

11.12.4 Automotive heterogeneous AI systems

Automotive applications combine constrained power and cooling with deadlines and functional-safety requirements. The required guarantees depend on the function's safety classification, so not every in-vehicle ML workload is hard real time.

Automotive SoCs use combinations of redundancy, monitoring, partitioning, and bounded scheduling to support safety objectives while accelerators execute perception and other compute-intensive functions. Shared bandwidth is one concrete risk: a convenience workload can interfere with a perception pipeline that has a validated deadline. Hardware partitions, bandwidth controls, or time-triggered schedules can isolate the relevant resources, but the mechanism and guarantee are platform-specific. Multisensor pipelines likewise need bounded end-to-end timing rather than high average throughput alone.

Vehicles may distribute sensing and computation across several controllers or consolidate functions onto central platforms. Either design must timestamp, transport, and align sensor data within the

assumptions of the fusion pipeline. The resulting synchronization and communication costs belong in the same latency budget as accelerator execution.

External vehicle-to-everything (V2X) inputs add another source with different latency and trust properties. A safe design cannot assume those messages arrive with the timing or reliability of local sensors; any use in a safety-related function must be bounded and validated accordingly.

11.12.5 Software stack challenges

The architectural sophistication of heterogeneous SoCs turns software into the coordination layer for power, thermal state, determinism, and operator fallback. OpenCL targets heterogeneous compute devices, while Vulkan exposes portable GPU compute; neither erases device-specific tuning. Mobile inference runtimes use delegates or backends for CPUs, GPUs, DSPs, and NPUs, but unsupported operators may fall back to another processor and add transfers. Deployment therefore requires inspecting the realized partition, not only the framework graph.

Shared memory architectures compound the coordination problem. Memory management must account for processor-specific caching behavior, memory access patterns, and coherency requirements. CPU caches may interfere with GPU memory access patterns, while NPU direct memory access (DMA) operations must be synchronized with CPU cache operations to maintain data consistency.

Heterogeneous SoCs address this complexity with profiled cost models, rules, control policies, and sometimes learned predictors. Telemetry on latency, utilization, temperature, and power can guide selection among validated execution plans; it does not remove the need to characterize operator support and worst-case behavior.

No single processor architecture can optimally handle the diverse computational patterns in AI applications, so heterogeneous acceleration becomes a coordination problem rather than a hardware inventory. Efficient mobile AI systems deliver high performance only when processor assignment, memory coherence, thermal limits, and latency constraints are managed together.

The same coordination problem also has an energy consequence. If hardware selection determines how much data moves, how often accelerators stall, and how efficiently arithmetic maps to silicon, then it also determines how much energy the deployment consumes for each useful prediction.

11.13 Hardware Sustainability

At fleet scale, energy per useful inference becomes an important hardware-selection criterion. Operational emissions depend on workload throughput, measured system power, utilization, cooling overhead, and the electricity mix; embodied emissions add another lifecycle cost. Peak performance per watt is therefore a screening metric, not a substitute for end-to-end energy per request. The following scenario shows how assumptions compound when comparing a generic-CPU fleet with specialized accelerators.

Napkin Math 11.10: An illustrative operational-energy comparison

Problem: Under a stated throughput-equivalence assumption, how does operational energy change when a CPU fleet is replaced by specialized NPUs (Neural Processing Units)?

Assumptions: This stylized comparison treats the listed peak rates as useful workload throughput, assumes both fleets run continuously at the listed power, and excludes host power, cooling, embodied carbon, and utilization differences. It is a sensitivity calculation, not a measured procurement result.

- **CPU inference:** 100 W for 1 TFLOP/s (efficiency = 0.01 TFLOP/s/W).
- **NPU inference:** 5 W for 10 TFLOP/s (efficiency = 2 TFLOP/s/W).
- **The assumed peak-efficiency gap:** 200×

Math:

1. **Workload:** 1 billion inferences per day.
2. **CPU fleet energy:** 1,000 CPU servers $\times$ 100 W $\times$ 24 h $\approx$ 2,400 kWh/day.

3. **NPU fleet energy:** $100 \text{ NPU chips} \times 5 \text{ W} \times 24 \text{ h} \approx 12 \text{ kWh/day}$.
4. **Carbon savings:** At 0.429 kg/kWh , switching to NPUs saves $\sim 373.9 \text{ t}$ of CO_2 per year.

Systems insight: Specialized accelerators can reduce operational energy when the workload uses their efficient path and fewer devices deliver the required service throughput. Replace these assumptions with measured requests per second and wall power before using the result for procurement or carbon accounting.

The sustainability perspective reinforces a theme that has recurred throughout this chapter: hardware selection is not determined by peak speed alone. Energy per useful result, carbon intensity, embodied impact, and total cost of ownership belong beside latency, throughput, and memory capacity. The remaining step is to identify the misconceptions that cause teams to choose the wrong hardware path.

11.14 Fallacies and Pitfalls

Hardware acceleration has counterintuitive performance characteristics because impressive specifications can mask workload-specific bottlenecks. The fallacies and pitfalls here capture selection and optimization errors that leave expensive resources underused.

Fallacy: *More specialized hardware always provides better performance than general-purpose alternatives.*

Engineers assume specialized accelerators automatically outperform general-purpose processors for all AI workloads. In reality, specialized hardware achieves peak performance only when workloads match architectural assumptions, the core of algorithm-machine co-design. As demonstrated in section 11.6, operations must exceed the accelerator's ridge point to be compute bound; an A100 GPU has a ridge point of 153 FLOP/byte, meaning operations with arithmetic intensity below this threshold are memory bound regardless of the accelerator's 312 TFLOP/s peak compute. A transformer attention softmax with $\text{AI} = 2 \text{ FLOP/byte}$ – 5 FLOP/byte achieves only 4.1 TFLOP/s – 10.2 TFLOP/s (3.3 percent utilization) on an A100. CPUs with ridge points around 10 FLOP/byte – 20 FLOP/byte still treat this kernel as memory-bound, but the same AI range corresponds to about 10 percent–50 percent of a CPU's lower peak. Models with irregular memory access, small batch sizes, or dynamic computation graphs may perform better on flexible processors. Effective hardware selection requires matching workload arithmetic intensity to architectural ridge points, not assuming specialization always wins.

Pitfall: *Ignoring memory bandwidth limitations when selecting acceleration strategies.*

Practitioners focus on peak TFLOP/s without analyzing whether their workloads can achieve compute-bound performance. As quantified in section 11.5.1, the energy model used here assigns about 640 pJ to a DRAM access versus 0.5 pJ for an on-chip L1 SRAM access, creating orders-of-magnitude energy penalties. An accelerator advertising 300 TFLOP/s with 2 TB/s bandwidth has a ridge point of 150 FLOP/byte ; LayerNorm operations with $\text{AI} = 1.5 \text{ FLOP/byte}$ achieve only 3 TFLOP/s (1 percent utilization) in this worked example. Organizations can deploy expensive high-compute accelerators for memory-bound workloads and still see low utilization if bandwidth, not compute, is the bottleneck. Teams must calculate workload arithmetic intensity and compare against hardware ridge points before purchasing accelerators.

Fallacy: *Hardware acceleration benefits scale linearly with additional accelerators.*

Teams expect eight GPUs to train $8\times$ faster than one GPU. Multi-accelerator scaling introduces communication overhead that violates linear scaling assumptions. As noted in section 11.11, AllReduce operations for gradient synchronization can require exchanging large gradient payloads for large models. In an ideal ring AllReduce, each rank transfers $2(N-1)/N$ payloads, or $1.75\times$ payloads for eight GPUs. With NVLink delivering 600 GB/s bidirectional (half that per direction), synchronizing a 1 GB gradient therefore requires 5.83 ms ; relative to a 50 ms compute step, this is 11.7 percent before protocol overhead. Without compute-communication overlap, this worked eight-GPU scenario achieves about $7.2\times$ speedup (89.6 percent efficiency) before load imbalance, synchronization barriers, and insufficient parallel work reduce scaling further.

Pitfall: *Planning accelerator capacity from peak FLOP/s specifications.*

Peak FLOP/s is one accelerator capability, but delivered throughput also depends on arithmetic intensity, shape, precision, kernel efficiency, and system overhead. The Roofline Model (section 11.6) bounds the compute and bandwidth portions of that gap. In the representative budgeting scenarios here, A100 transformer training sustains 120 TFLOP/s–180 TFLOP/s, while a sparse recommender workload sustains 10 TFLOP/s–30 TFLOP/s. These ranges are planning assumptions, not universal measurements. Capacity plans should use representative measured throughput where possible and roofline estimates when measurements are unavailable.

Fallacy: *Any FLOP/s rating can estimate a low-precision workload.*

Accelerators have separate datapaths for different precisions, and the peak throughput varies dramatically across them. An H100 delivers roughly 989 TFLOP/s in FP16 tensor operations but only about 67 TFLOP/s in FP32 CUDA-core operations: a roughly 15× gap within the same chip (Choquette 2023). Estimating training time with the FP32 number when the workload actually uses BF16 produces utilization figures that look catastrophic for no reason, and matching against the wrong roofline misclassifies kernels as compute-bound when they are memory-bound (or vice versa). Always match the peak constant to the precision the workload actually issues, and quote precision explicitly when reporting model FLOPs utilization.

Pitfall: *Deploying small-batch inference workloads on high-compute accelerators.*

Small-batch inference can underuse a high-compute accelerator. In this dense-layer traffic model, $M = N = 2048$ gives $AI = 1$ FLOP/byte at batch 1 and $AI = 204.8$ FLOP/byte at batch 256. At batch 1, the memory-bound roofline ceilings are about 2.04 TFLOP/s on A100 and 0.3 TFLOP/s on T4. The T4’s FP16 Tensor Core peak is 65 TFLOP/s, with a ridge point of 203.1 FLOP/byte. A lower-cost accelerator can therefore be more economical for this regime, but latency, memory capacity, concurrency, software support, and current pricing must be measured. Match the deployment to its actual batch and service objective rather than assigning accelerator classes by name.

Fallacy: *Vendor-specific optimizations have no long-term portability cost.*

Organizations optimize exclusively for specific vendors to maximize performance without considering system flexibility. As discussed in section 11.9, deep integration with vendor-specific libraries (CUDA, TensorRT, XLA) and custom kernels creates lock-in. A codebase with many hand-written accelerator kernels can require substantial engineering effort to port to a different vendor, delaying hardware upgrades and preventing multi-vendor deployments. Vendor-specific optimizations should therefore be isolated behind hardware abstraction layers. Maintaining portable code paths enables vendor competition, hardware flexibility, and faster adoption of emerging accelerators while still capturing most performance benefits through framework-level optimizations.



Checkpoint 11.4: Feasibility assessment: Can you run it?

Before procuring hardware, validate all three hard constraints. The fallacies and pitfalls in this section reduce to a concrete procurement test:

- ☐ **Memory capacity:** compute $M_{\text{req}} = \text{Weights} + \text{KV Cache} + \text{Activation Buffer}$ and verify $M_{\text{req}} < M_{\text{device}}$ for a 7-billion-parameter Llama model with 14 GB FP16 weights on a 16 GB GPU.
- ☐ **Bandwidth:** compute $T_{\text{token}} = D_{\text{vol}}/\text{BW}$ for a 70-billion-parameter model (140 GB) on 1 TB/s memory bandwidth, then compare the result with a 50 ms latency target.
- ☐ **Compute:** compute $T_{\text{process}} = O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$ and compare it with the throughput target. Processing video at 30 FPS requires completing inference within 33.3 ms.
- ☐ **Roofline placement:** for an accelerator at roughly 2 TFLOP/s peak and 200 GB/s bandwidth, estimate the ridge-point arithmetic intensity, then place a 10 FLOP/byte kernel against it: is it compute- or memory-bound, and would a faster clock help it at all?

This checklist synthesizes the principles developed throughout this chapter, translating theoretical understanding into practical engineering decisions. Together, these fallacies reduce the chapter’s machinery to a diagnostic habit of starting from the workload, choosing the bottleneck metric, matching the hardware path, and budgeting for its portability, scaling, and energy consequences.

11.15 Summary

Hardware acceleration co-designs compute engines, memory hierarchies, numerical formats, and dataflows. For a chosen memory interface, the iron law compares compute time $O/(R_{\text{peak}}\eta_{\text{hw}})$ with movement time D_{vol}/BW ; the roofline restates that competition as $\min(R_{\text{peak}}, \text{BW}I)$. End-to-end latency also includes launch, synchronization, communication, and fixed overheads. Accelerators improve efficiency only when supported paths capture reuse locally. Hardware selection therefore begins with the workload’s measured or modeled bottleneck, not a vendor peak.



Key Takeaways: Moving data costs more than computing it

- **The Roofline Model identifies bottlenecks:** Compare a kernel’s measured intensity with the target’s precision-specific ridge point (153 FLOP/byte here). Below it, reduce memory traffic; above it, improve compute utilization.
- **Memory bandwidth constrains performance:** GPU compute capacity has grown faster than memory bandwidth. Many inference phases are memory bound, especially low-batch decoding and embedding lookup, while large-batch GEMMs can remain compute bound.
- **Hardware-software co-design compounds performance:** Matching algorithm patterns to architectural capabilities, such as dense GEMM on systolic or tensor units and supported sparsity on sparse paths, can produce large improvements.
- **Tensor Cores require a compatible path:** Precision, dimensions, layouts, and library support must match the architecture. Batch and shape affect reuse and utilization, but no single batch threshold guarantees peak performance.
- **End-to-end speedup stops at the next bottleneck:** Accelerating one kernel leaves launch, transfer, synchronization, communication, and unsupported work unchanged. Use Amdahl’s law and profiling to decide whether further silicon, better mapping, or less movement will improve the complete workload.

An accelerator is not a uniformly faster computer; it provides efficient paths for particular operations and dataflows. Tiling, fusion, hierarchy-aware scheduling, and systolic execution reduce modeled-interface traffic only when the workload, numerical format, and software path can use them.


The Roofline Model provides a first diagnostic by using arithmetic intensity and the ridge point to narrow the likely bottleneck; profiling then accounts for latency, occupancy, communication, and implementation overhead. The result is a testable hypothesis, not a prediction.



What’s Next: From optimization to validation

The ML optimization stack rests on three core efficiency pillars: data selection (chapter 9) minimizes training sample requirements, model compression (chapter 10) reduces algorithmic parameter and precision complexity, and hardware acceleration (developed in this chapter) maximizes physical machine throughput. Optimization without empirical measurement, however, remains incomplete. Chapter 12 moves from theoretical FLOP bounds to measured wall-clock latency, applying standardized benchmark suites and statistical protocols to test optimization claims against operational reality.

Benchmarking




12.1	ML Benchmarking Framework
12.2	Historical Foundations
12.3	System Benchmarking Suites
12.4	Benchmarking Granularity
12.5	Benchmark Components
12.6	Training vs. Inference
12.7	Training Benchmarks
12.8	Inference Benchmarks
12.9	Power Measurement Techniques
12.10	Benchmarking Best Practices
12.11	Model and Data Evaluation
12.12	Production Considerations
12.13	Fallacies and Pitfalls
12.14	Summary

Purpose

How can ML systems be compared fairly when hardware, models, data, and deployment all interact?

Benchmarking brings together decisions already developed through data selection, model compression, and hardware acceleration, then tests whether their gains survive under deployment-representative conditions. Each decision targets one dimension (latency, accuracy, throughput, or energy), but an ML system is the product of all these dimensions simultaneously. A pruned model runs faster on one accelerator but slower on another. A larger batch size improves accelerator utilization but can violate a latency service-level agreement. An edge device may advertise peak throughput that thermal throttling sharply reduces under sustained workloads. The challenge is not whether a local metric improves but whether the combined system improves under conditions that actually matter. Benchmarking makes such comparisons systematic rather than anecdotal. It requires defining what to measure (accuracy, latency, throughput, energy), at what granularity (a single kernel, a full model, an end-to-end pipeline), and under which conditions (batch size, input distribution, thermal state, concurrent load). Without this structure, teams compare numbers that were never measured on the same terms, and decisions that looked sound in a spreadsheet collapse under production workloads. Earlier chapters optimized the model, selected the data, and matched the hardware. Benchmarking validates those optimizations, bringing claims into contact with evidence and quantifying the gap between promise and delivery. In D·A·M terms, benchmarking holds co-design to account by revealing whether Data, Algorithm, and Machine were matched or merely assembled.





Learning Objectives

- Explain benchmarking as D·A·M validation that tests whether optimization claims hold under representative conditions
- Compare training and inference benchmarks using throughput, latency percentiles, energy, accuracy, and workload scope
- Select micro, macro, or end-to-end granularity based on the engineering decision being tested
- Apply standardized benchmark run rules to align datasets, metrics, hardware configuration, and reporting
- Design benchmark protocols that control power boundaries, input distributions, batch sizes, and statistical variance
- Evaluate model and data quality with calibration, robustness, representativeness, and slice-level metrics
- Diagnose benchmark-production gaps caused by drift, thermal throttling, dynamic load, and silent degradation

12.1 ML Benchmarking Framework

A model quantized to INT8 may benchmark $2\times$ faster on a synthetic workload but show no improvement under real traffic patterns with variable input sizes and concurrent requests. A pruned model may maintain accuracy on the test set but fail on edge cases the benchmark never covered. Data selection promises more efficient training, model compression promises smaller and faster models, and hardware acceleration promises higher throughput. Verifying that these claims hold in production is itself an engineering discipline.



Definition 12.1: Machine learning benchmarking

Machine learning benchmarking is the empirical measurement of ML workloads under specified conditions, used to test claims about components, models, or end-to-end systems against representative evidence rather than peak specifications alone.

1. **Significance:** The gap between peak and sustained performance can be large. An A100 GPU delivers 312 TFLOP/s (BF16) at peak, but in this illustrative 30 percent–50 percent MFU scenario it sustains 93.6 TFLOP/s–156 TFLOP/s, about a $2\text{--}3.3\times$ gap due to factors such as memory stalls, pipeline bubbles, and kernel launch overhead. Benchmarking quantifies the actual η_{hw} gap for a workload; vendor spec sheets do not.
2. **Distinction:** ML benchmarking spans multiple scopes. Micro-benchmarks isolate operations such as matrix multiplication, model-level benchmarks evaluate complete training or inference workloads, and end-to-end benchmarks include surrounding work such as data loading, preprocessing, optimization, checkpoint I/O, or serving infrastructure. The scope must match the engineering claim.
3. **Common pitfall:** A frequent misconception is that benchmark numbers are stable references. Both the workload (new model architectures) and the hardware (new GPU generations) evolve, so a result that leads a benchmark under one version often becomes the baseline under a later version, making year-over-year comparisons meaningful only when the benchmark version is held constant.

Benchmarking is where the physical laws established in earlier chapters face empirical reality. The **benchmark-production gap** measures the difference between modeled expectations and observed production behavior. Closing that gap requires measurements that convert theoretical claims into verified engineering knowledge.

ML benchmarking operates across three interdependent dimensions that map directly to the components of any deployed system. **System benchmarking** measures whether the hardware delivers

promised performance under realistic workloads or whether memory bandwidth saturation and software dispatch overhead erode the gains. **Model benchmarking** measures whether optimization techniques preserve model quality across the full input distribution, not just on curated test sets. **Data benchmarking** measures whether the model generalizes to real-world data with all its noise, bias, and distributional shift. Each dimension can independently reveal problems invisible to the others, and a system that passes all three provides far stronger deployment confidence than one evaluated along any single axis.

An ML benchmark captures a snapshot of a workload and data distribution rather than a permanent specification. The gap between peak and sustained performance is not fixed either; it shifts as workloads and hardware generations evolve, making any single benchmark result time-stamped rather than universal.



Systems Perspective 12.1: Benchmarks as moving targets

In traditional systems (for example, SPEC CPU), the benchmark is a *rigid specification*. A sorting algorithm is correct if it sorts the list. Correctness is absolute and unchanging. In ML systems, the benchmark is a *soft specification*: correctness is defined by a finite set of examples (ImageNet), and the world moves. A model that scores highly on ImageNet can still underperform on user photos taken years after the benchmark was created.

In computer architecture, engineers design for the benchmark because the benchmark represents the workload. In ML engineering, designing solely for the benchmark is *overfitting*. Robustness comes from acknowledging that the benchmark is only a proxy for a shifting reality.

MobileNetV2 deployment validation makes the three-dimensional framework concrete. It serves as the chapter's running example because it spans all three evaluation dimensions, illustrating how each reveals problems the others cannot.



Lighthouse 12.1: MobileNetV2 deployment validation

MobileNetV2 (introduced in section 6.1.1) serves as the lighthouse example for validating the complete optimization pipeline. MobileNetV2 refines v1's depthwise separable design with inverted residuals and linear bottlenecks while maintaining a similar parameter scale (Sandler et al. 2018). It exemplifies the deployment challenges where benchmarking determines success or failure: compression can reduce model size, hardware acceleration can reduce inference latency, and only benchmarking can determine whether those gains survive the full deployment pipeline.


1. **Model compression** (chapter 10): INT8 quantization reduces this MobileNetV2 worked example from 14 MB to 3.5 MB (4× compression)
2. **Hardware acceleration** (chapter 11): the illustrative EdgeTPU scenario uses 2 ms inference vs. 15 ms on CPU
3. Benchmarking validation: Verify the pipeline delivers in practice

The sections that follow address one dimension of this validation stack at a time, building toward a systematic methodology that isolates EdgeTPU latency from preprocessing and data transfer overhead, confirms INT8 quantization preserves accuracy on edge cases such as unusual lighting, and checks that performance holds on real-world smartphone images rather than only ImageNet test images.

Rigorous evaluation begins with the mindset that separates meaningful evidence from misleading metrics. Three principles distinguish effective practitioners.

First, *benchmarks are proxies, not truth*. Every benchmark measures specific conditions that may not match the target deployment. A system can achieve high sample throughput in Offline mode (bulk throughput with all inputs available) and much lower queries per second (QPS) in Server

¹ | **Goodhart's law:** Goodhart (1984) articulated the original 1975 Bank of England observation on monetary policy; Strathern (1997) generalized it into the form quoted above. The original context was macroeconomics: once a monetary aggregate became an official policy target, banks changed behavior to game the metric, destroying its predictive value. In ML, the same failure mode recurs structurally: BLEU rewards n-gram overlap (Papineni et al. 2002), ImageNet rewards performance on a fixed visual distribution (Deng et al. 2009; Recht et al. 2019), and benchmark leaderboards can incentivize test-set-specific tuning.



Component speedup rarely survives as end-to-end benchmark speedup.

² | **Benchmark:** From surveying, where a “bench mark” was a horizontal cut in stone serving as a fixed elevation reference. The term entered computing in the 1970s to describe standardized comparison points, but the surveying metaphor carries a systems lesson: just as an elevation measurement is meaningless without a calibrated reference, an ML throughput number is meaningless without controlled workloads, thermal state, and precision settings.

mode (latency-constrained requests arriving over time). The critical question is always what the benchmark does not measure.

Second, Goodhart's Law applies everywhere.¹ “When a measure becomes a target, it ceases to be a good measure.” Teams that optimize for benchmark rankings often produce systems that excel in evaluation but fail in production. Benchmark-specific optimizations frequently degrade characteristics that matter for deployment: robustness, calibration, and efficiency.

Third, end-to-end beats component metrics. In this illustrative pipeline, a 3× inference speedup applied to a 10 ms model stage inside a 50 ms request yields only about a 1.2× end-to-end improvement, or worse if the optimization increases memory pressure. These principles reappear throughout the benchmarking methodology and are examined in depth in section 12.13.

Knowing *what* to measure, however, is only half the problem. Measuring incorrectly (with the wrong workloads, biased baselines, or uncontrolled variables) produces numbers that feel precise but mislead decisions. The history of computing benchmarking is littered with examples of technically sound metrics applied with flawed methodology, from compiler-gamed Whetstone scores to cherry-picked GPU benchmarks that predict nothing about sustained workloads. Understanding how measurement methodology evolved, and where it failed, is essential for designing benchmarks that distinguish genuine improvements from measurement artifacts.

The historical foundations of benchmarking² matter because they expose the validation failures that still recur in ML: optimized metrics that stop predicting real workloads, hardware numbers that ignore sustained operating state, and model scores that miss deployment cost. The same validation sequence governs modern practice: first verify that hardware delivers promised performance, then verify that the model and data optimizations built atop that hardware deliver their promised gains.

12.2 Historical Foundations

In 1976, when Whetstone became one of the first standardized computing benchmarks, vendors began optimizing their compilers specifically for its floating-point tests, producing impressive numbers that did not reliably predict real application performance. Similar gaming has affected later benchmark generations. Understanding *why* ML benchmarking requires a three-dimensional approach demands tracing *how* measurement methodologies evolved, and often failed, over decades of computing history. Each generation of benchmarks emerged from the limitations of its predecessors, teaching lessons that directly inform modern ML evaluation.

Before that history begins, one boundary condition matters: a benchmark is useful only when it names the layer whose claim it validates.

🔍 Systems Perspective 12.2: Benchmarking as cross-layer evidence

An ML benchmark is meaningful only when it identifies which layer of the system is being validated. Data selection metrics such as PPD, area under the learning curve (AULC), and data compression ratio (DCR), all developed in chapter 9, measure whether fewer examples can preserve learning signal. Compression metrics (chapter 10) measure whether fewer parameters or lower precision preserve quality. Hardware metrics such as roofline position and TOPS/W (chapter 11) measure whether the machine executes the workload efficiently. The system benchmarks in section 12.3.2 through section 12.9.4 connect these layers by testing whether local improvements survive in an end-to-end workload.

That cross-layer role explains why benchmark history matters: each generation of performance measurement advanced when practitioners discovered that the previous method failed to predict real-world behavior. The evolution from simple performance metrics to ML benchmarking reveals three methodological shifts.

12.2.1 Performance benchmarks

The earliest computing benchmarks revealed a problem that plagues evaluation to this day: benchmark gaming. Whetstone (Curnow and Wichmann 1976) used a synthetic mix of scientific-program operations, while LINPACK³ (Dongarra et al. 1979) measured dense linear-system solving. Vendors

could optimize specifically for such fixed tests rather than for broader workloads. SPEC CPU (1989) broadened evaluation through a suite of portable, application-oriented programs (Dixit 1993). This lesson directly shapes ML benchmarking: optimization claims from chapter 10 require validation on representative tasks, and MLPerf’s inclusion of models such as ResNet-50 and BERT captures more of the deployment stack than an isolated kernel.

As deployment contexts diversified, a second limitation emerged: single-metric evaluation proved inadequate. Graphics benchmarks began measuring rendering quality alongside frame rate; mobile benchmarks added battery life as a co-equal concern with performance. The multi-objective challenges from chapter 1 (balancing accuracy, latency, and energy) manifest directly in ML evaluation, where no single metric captures deployment viability.

A third shift occurred when distributed computing revealed that component-level optimization fails to predict system-level performance. A CPU benchmark cannot predict cluster throughput when network communication dominates. ML training similarly depends on the interplay of accelerator compute (chapter 11), data pipelines, gradient synchronization, and storage throughput. MLPerf evaluates complete workflows, recognizing that performance emerges from component interactions, not from components in isolation.

DAWNBench (Coleman et al. 2019) emerged as an early ML benchmark that pioneered time-to-accuracy evaluation, directly influencing MLPerf’s methodology for measuring training efficiency. These lessons culminate in MLPerf⁴ (2018), which synthesizes representative workloads, multi-objective evaluation, and integrated measurement while addressing ML-specific challenges (Mattson et al. 2020; Reddi et al. 2019).

12.2.2 Energy benchmarks

The multi-objective evaluation paradigm naturally extended to energy efficiency as computing diversified beyond mainframes with less constrained power budgets. Mobile devices demanded battery life optimization, while warehouse-scale systems faced energy costs rivaling hardware expenses. This shift established energy as a first-class metric alongside performance, spawning benchmarks like SPEC Power⁵ for servers and Green500⁶ for supercomputers.

Diverse workload patterns and system configurations continue to challenge power benchmarking across computing environments. MLPerf Power (MLCommons 2024b) addresses this with specialized methodologies for measuring the energy impact of machine learning workloads, reflecting energy efficiency’s central role in AI system design.

Energy benchmarking extends beyond hardware power measurement to include algorithmic efficiency. Model compression techniques (pruning, quantization, knowledge distillation) can reduce energy by changing the work a system performs, not only by changing the hardware that performs it. MobileNet-family architectures use depthwise separable convolutions to cut computation relative to heavier convolutional neural network (CNN) baselines such as ResNet (Howard et al. 2017; Sandler et al. 2018; He et al. 2016a). These techniques, detailed in chapter 10, establish that energy-aware benchmarking must evaluate algorithmic efficiency alongside hardware power consumption; section 10.4.1.1 quantifies the specific energy breakdown of INT8 vs. FP32. As AI systems scale, this lesson becomes central to sustainable computing practices.

12.2.3 Domain-specific benchmarks

As computing diversified beyond general-purpose servers, generic benchmarks proved inadequate for specialized domains. Three categories of specialization drove this evolution, each exposing measurement dimensions that general-purpose benchmarks could not address.

Deployment constraints shape core metric priorities. Data center workloads optimize for throughput with rack- and cluster-scale power budgets, while mobile AI operates within tight device thermal envelopes, and IoT devices require milliwatt-scale operation. These constraints, rooted in efficiency principles from chapter 1, determine whether benchmarks prioritize total throughput or energy per operation.

Application requirements then impose functional and regulatory constraints beyond raw performance. Healthcare AI demands interpretability metrics alongside accuracy; financial systems may require very low latency with audit compliance; autonomous vehicles need safety-critical reliability and formal functional-safety validation. These requirements extend evaluation beyond traditional

³ | **Whetstone and LINPACK:** Whetstone (Curnow and Wichmann 1976) was named after the English Electric facility in Whetstone, Leicestershire, where the original ALGOL compiler was built; LINPACK (Dongarra et al. 1979) was a package and benchmark for dense linear systems, later used by the Top500 list. Whetstone’s fixed synthetic program mix and LINPACK’s dense linear-algebra focus made each useful but narrower than a diverse application suite. ML benchmarking inherited the same vulnerability: model-specific kernel tuning can overfit a single workload, which is why MLPerf uses multiple workloads spanning vision, language, and recommendation (Mattson et al. 2020; Reddi et al. 2019).

⁴ | **MLPerf:** Launched in 2018 by a consortium of industry and academic institutions, MLPerf takes its name from “ML” combined with “Perf” (performance), echoing SPEC’s benchmarking tradition. MLPerf’s design principles—representative workloads, full-system measurement, and open submission—directly address the gaming that plagued Whetstone and LINPACK: vendors who could previously report peak kernel throughput on cherry-picked problem sizes must now report end-to-end system performance on standardized tasks (Mattson et al. 2020; Reddi et al. 2019).

⁵ | **SPEC Power:** Introduced in 2007, SPEC Power measures performance per watt across 11 load levels from idle (0 percent) through 100 percent in 10 percent increments (Lange 2009). This granularity matters for ML serving: inference workloads rarely sustain 100 percent load, and servers that are efficient at peak but wasteful at partial load inflate the energy cost of real-world deployment.

⁶ | **Green500:** Started in 2007 as a counterpart to the Top500, Green500 ranks systems by FLOP/s per watt rather than raw performance (Feng and Cameron 2007). Its lesson for ML systems is methodological: the most cost-effective training cluster is not necessarily the fastest one, but the system that delivers useful work per watt under the workload and measurement boundary that matter.

performance metrics; chapter 15 later systematizes the responsible-engineering principles behind fairness, interpretability, and compliance.

Operational conditions determine real-world viability. Autonomous vehicles face wide temperature ranges and degraded sensor inputs; data centers handle large concurrent request volumes with network faults; industrial IoT endures long deployments without maintenance. The hardware capabilities from chapter 11 only deliver value when validated under these conditions.

Machine learning exemplifies this transition to domain-specific evaluation. Traditional CPU and GPU benchmarks prove insufficient for assessing ML workloads, which involve complex interactions between computation, memory bandwidth, and data movement patterns. MLPerf provides standardized performance measurement for machine learning models across these categories: MLPerf Training addresses data center deployment constraints with multi-node scaling benchmarks (Mattson et al. 2020), MLPerf Inference evaluates latency-critical application requirements across server to edge deployments (Reddi et al. 2019), MLPerf Tiny assesses ultra-constrained operational conditions for microcontroller deployments (C. Banbury et al. 2021), and a cross-cutting MLPerf Power track measures energy efficiency under each of these regimes. Reading table 12.1 down its constraint column shows tighter limits as deployment scale shrinks: data-center interconnect bandwidth gives way to latency service level agreements (SLAs) at server and edge, then to ultra-low-power operation with kilobytes of memory on microcontrollers. The same three-category framework, applied to each scale, produces a suite whose metrics track what actually limits the system at that scale rather than a single universal score.

Table 12.1: MLPerf Benchmark Suite Variants: Each variant addresses a different deployment context, from data-center-scale training to ultra-constrained microcontroller inference, targeting specific operational constraints and measuring metrics relevant to its deployment scenario.

MLPerf Variant	Target Domain	Key Constraints	Primary Metrics
MLPerf Training	Data center	Multi-node scaling, high bandwidth interconnects	Time-to-quality, throughput (samples/sec)
MLPerf Inference	Server/Edge	Latency SLAs, throughput requirements	QPS, latency percentiles, accuracy preservation
MLPerf Tiny	MCU/IoT	Ultra-low-power inference, limited memory	Latency, accuracy, energy per inference
MLPerf Power	Cross-cutting	Energy budgets, thermal constraints	Performance/W, energy per query

MLPerf Power extends the same discipline to energy efficiency, where the benchmarked quantity is useful work per watt rather than raw throughput alone. Domain-specific benchmarks drive targeted hardware and software optimizations while ensuring that improvements translate to deployment success rather than narrow laboratory conditions.

This historical progression, from general computing benchmarks through energy-aware measurement to domain-specific evaluation frameworks, provides the foundation for understanding ML benchmarking challenges. The lessons learned (representative workloads over synthetic tests, multi-objective over single metrics, integrated systems over isolated components) directly shape AI system evaluation. Table 12.2 summarizes this progression and the key lessons each generation contributed.

These lessons culminate in ML benchmarking suites, yet ML systems add statistical and data-dependent variability to the sources of system noise found in other workloads. They must satisfy all three historical lessons (representative workloads, multi-objective evaluation, integrated measurement) while also accounting for outcomes that vary with training data, weight initialization, and operation ordering. This additional variability requires corresponding statistical controls.

Individual organizations learned these lessons independently, often painfully, but isolated measurements cannot drive an industry. When one team measures inference latency including preprocessing and another excludes it, when accuracy benchmarks use different data splits, or when power measurements draw different system boundaries, the resulting numbers are incommensurable. The transition from ad-hoc measurement to standardized benchmarking suites transforms benchmarking from an internal validation exercise into a shared language that enables hardware procurement, architecture comparison, and deployment decisions across organizations.

Table 12.2: Benchmark Evolution: Evolution of computing benchmarks from synthetic operations to ML-specific evaluation. Each generation addressed limitations of its predecessors, culminating in MLPerf’s synthesis of representative workloads, multi-objective metrics, and integrated system measurement.

Benchmark	Year	Primary Focus	Key Metric(s)	Lesson for ML Benchmarking
Whetstone	1976	Synthetic floating-point operations	MWIPS	Gaming synthetic tests undermines evaluation validity
LINPACK	1979	Linear algebra (matrix operations)	FLOP/s	Isolated operations miss system-level complexity and bottlenecks
SPEC CPU	1989	Real application workloads	SPECrate, SPECspeed	Representative workloads reveal true deployment performance
SPEC Power	2007	Server energy efficiency	ssj_ops/W across load levels	Energy efficiency requires multi-load evaluation, not just peak performance
Green500	2007	HPC energy efficiency	GFLOP/s per watt	Efficiency rankings complement raw performance rankings
MLPerf	2018	ML systems (training + inference)	Time-to-quality, QPS, latency, accuracy	Synthesizes all lessons: representative workloads + multi-objective + system

12.3 System Benchmarking Suites

A team evaluating edge deployment hardware needs to compare five different system on chip (SoC) designs for a smart camera product. Vendor A reports 8 TOPS at INT8; Vendor B reports 15 TOPS at INT4; Vendor C reports inference latency on a proprietary model; Vendor D cites MLPerf scores from two generations ago; Vendor E provides only peak throughput at maximum batch size. None of these numbers are comparable. The team cannot make a procurement decision because every vendor measured a different thing, under different conditions, using different definitions of “performance.” The problem is not *a lack of data* but *a lack of commensurable data*, and benchmarking suites exist to solve exactly this fragmentation.

Three lessons from benchmark history (representative workloads, multi-objective evaluation, and integrated measurement) converge with the challenge unique to ML: inherent probabilistic variability. Modern benchmarking suites encode these lessons into standardized frameworks that make the kind of cross-organization comparison the hardware procurement team needs possible.

ML benchmarks must evaluate the interplay between algorithms, hardware, and data, not merely computational efficiency alone. Early benchmarks focused on algorithmic performance (LeCun et al. 1998), but scaling demands expanded the focus to hardware efficiency (Jouppi et al. 2017), and high-profile deployment failures elevated data quality as a third evaluation dimension (Gebu et al. 2021). This probabilistic nature elevates accuracy to a first-class evaluation dimension alongside speed and energy consumption: the same ML system can produce different results depending on the data it encounters. Energy efficiency cuts across all three framework dimensions, since algorithmic choices affect computational complexity (Hernandez and Brown 2020), hardware capabilities determine energy-performance trade-offs, and dataset characteristics influence training energy costs.


12.3.1 ML measurement challenges

ML systems combine several sources of measurement variability that many traditional benchmarks were not designed to evaluate together. These include algorithmic randomness from weight initialization and data shuffling, hardware thermal states that affect clock speeds, system load variation from concurrent processes, and environmental factors such as network conditions and power management. Rigorous statistical methods are needed to distinguish genuine performance improvements from this measurement noise.

To address this variability, effective benchmark protocols require repeated experimental runs with random seeds chosen for the study design. The number of runs should follow the required uncertainty or statistical-power target, with measures beyond simple means (including standard deviations or confidence intervals) reported to quantify stability and distinguish genuine improvements from measurement noise.

Empirical studies have shown how inadequate statistical rigor can lead to misleading conclusions. Reinforcement-learning gains often fall within statistical noise (Henderson et al. 2018), while generative adversarial network comparisons often lack controlled experimental protocols, producing

inconsistent rankings across seeds (Lucic et al. 2018). These findings underscore the importance of establishing measurement protocols that account for ML’s probabilistic nature.



Napkin Math 12.1: The statistical confidence trap

Problem: A baseline classifier scores 95 percent, while a compressed version scores 94 percent on 1,000 images. Does this one-point drop establish a regression?

Math:

1. **Expected errors:** The scores correspond to 50 errors and 60 errors. Under the baseline binomial model, the error count has a standard deviation of about 7 errors.
2. **Difference interval (95 percent):** Let $\hat{p}_1$ and $\hat{p}_2$ denote the baseline and compressed accuracy estimates, respectively, and let N be the number of images evaluated for each model. Treating the two estimates as independent, the standard error of their difference is

$$SE(\hat{p}_1 - \hat{p}_2) = \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{N} + \frac{\hat{p}_2(1 - \hat{p}_2)}{N}}.$$

The observed 1 percentage point difference has an approximate interval from -1 percentage point to 3 percentage points, which includes zero.

Implication: Because the interval includes zero, the result does not establish a one-point regression. If both models ran on the same images, retain their disagreement counts and use a paired test such as McNemar’s test; marginal accuracies are insufficient. The 1,825-sample estimate for a single rate with ± 1 percentage point precision does not power this comparison.

Systems insight: Small benchmarks exhibit what amounts to a laboratory fallacy. The test set, viewed as a measurement instrument, must be sized to match the precision of the change it is meant to detect.

Representative workload selection determines benchmark validity. Synthetic microbenchmarks (which generate random uniform tensors in memory) often fail to capture the complexity of real ML workloads where data movement, memory allocation, and dynamic batching create performance patterns not visible in simplified tests. Conversely, **trace-based benchmarking** replays recorded production request traces—capturing realistic inter-arrival times, bursty arrival distributions, and variable sequence lengths—to stress-test serving infrastructure under true production conditions. Comprehensive benchmarking therefore requires workloads that reflect actual deployment patterns: variable sequence lengths in language models, mixed precision training regimes, and realistic data loading patterns that include preprocessing overhead.

Example 12.1: Goodhart’s Law in action

Scenario: An engineering team optimizes an NMT translation model for offline BLEU score, increasing beam search size from 1 to 10.

Diagnosis: The larger beam search achieves a 0.5-point BLEU gain on paper but increases candidate evaluations by 10 $\times$, causing inference latency to explode from 50 ms to 200 ms (4 $\times$ slowdown).

Systems lesson: Optimizing offline accuracy without latency constraints invites Goodhart’s Law failures. Production benchmarks must enforce strict service-level objectives (SLOs) for serving latency (e.g., latency < 100 ms).

Beyond workload representativeness, the distinction between *statistical significance* and *practical significance* requires careful interpretation. A small performance improvement might achieve statistical significance across hundreds of trials but prove operationally irrelevant if it falls within

measurement noise or costs exceed benefits. This is *the statistical confidence trap*: seemingly rigorous evaluation still misleads.

Statistical confidence is a measurement-capacity problem: the benchmark may be pointed at the right quantity, but the test set is too small to resolve the change. A second failure mode is metric alignment. Here the measurement can be precise and reproducible, yet still reward behavior that violates the deployed system's objective. The translation example makes that distinction concrete by showing how a BLEU improvement can come at the expense of latency.

These measurement failures share a deeper limitation: a benchmark on a static dataset measures recognition under a fixed distribution, not the robustness to a shifting one that production demands. The data dimension of the framework developed later in this chapter confronts exactly that gap.

These measurement challenges motivate evaluating each dimension of the three-dimensional framework (system, model, and data) with distinct methodologies. The bulk of this chapter focuses on system benchmarking (training benchmarks, inference benchmarks, and power measurement) because these form the foundation of standardized evaluation through MLPerf. Section 12.11 then develops the distinct methodologies required for model and data benchmarking.

12.3.2 System benchmarks

System benchmarks measure the computational foundation that enables model capabilities, examining how hardware architectures, memory systems, and interconnects affect overall performance. This validation is critical because hardware specifications often describe theoretical peaks that application workloads do not sustain. The discrepancy is common enough to make peak-performance claims incomplete. System benchmarks reveal these gaps by running standardized ML workloads rather than relying on peak arithmetic rates alone.

Systems Perspective 12.3: The fallacy of peak performance

Dave Patterson often refers to peak performance as “the performance the manufacturer guarantees you will not exceed.” For ML systems, the gap can be wide because of the memory wall: a memory-bound model leaves much of the advertised FLOP/s unreachable no matter how fast the arithmetic units run. Standardized benchmarks like MLPerf require systems to run specified models and data, revealing sustained performance at a documented operating point rather than a universal rate for every workload.

For memory-bound workloads, the peak-vs.-sustained gap follows from the memory wall; compute-bound workloads may approach peak. This distinction reframes vendor evaluation from guesswork into a checklist of concrete criteria.

Checkpoint 12.1: Decoding vendor benchmark claims

When evaluating hardware or software based on vendor-reported benchmarks, check whether the claim identifies the workload, measurement boundary, and operating conditions.

- ☐ Measured quantity: Does the latency number include preprocessing and postprocessing, not just model execution?
- ☐ Precision boundary: Can the precision behind a throughput claim, such as INT4 versus INT8 or FP16, be identified?
- ☐ Workload boundary: Are the model, batch size, and input distribution used to produce the comparison explicitly named?
- ☐ Excluded costs: Are memory transfers, model loading, initialization, sustained thermal behavior, and power accounted for at the claimed performance level?
- ☐ Claimed vs. achieved: Given a 1,000 TFLOP/s peak claim and a benchmark that sustains 350 TFLOP/s, estimate the MFU and state whether that utilization ratio alone identifies the bottleneck.

Engineers should reject any benchmark claim whose workload boundary, precision, and excluded costs cannot be reconstructed. A headline throughput or latency number becomes useful only after the engineer can map it to the actual model, batch shape, data movement, sustained operating point, and power envelope.

The underlying hardware—CPUs, GPUs, Tensor Processing Units (TPUs),⁷ and application-specific integrated circuits (ASICs)⁸—determines ML-system speed, efficiency, and scale. System benchmarks provide standardized methods for comparing hardware across AI workloads by computational throughput, memory bandwidth, power efficiency, operator coverage, and scaling (Reddi et al. 2019; Mattson et al. 2020).

Table 12.3 translates common marketing phrases into the technical caveats behind each.

Table 12.3: Decoding Vendor Benchmark Claims: Four common marketing phrases and the technical caveats that determine whether a claim describes end-to-end performance, a narrow accelerator measurement, or an unsupported efficiency comparison.

Vendor Claim	What It Often Means
“Up to 10,000 images/sec”	Peak throughput at maximum batch size, INT8, without preprocessing
“Sub-millisecond latency”	Accelerator compute only, excluding data transfer
“5× more efficient”	Per-operation efficiency, not total system efficiency
“Optimized for AI”	May only accelerate specific operations or precisions

System benchmarks serve two functions. For practitioners, they enable informed hardware selection by providing comparative data across configurations. For manufacturers, they quantify generational improvements and guide accelerator development. As GPU adoption grew, accuracy also improved rapidly, illustrating how hardware and algorithmic advances can drive progress together.

 Definition 12.2: Machine learning system benchmarks


Machine learning system benchmarks are standardized evaluation protocols that hold the workload and quality target constant while varying the hardware-software stack, measuring $\eta_{hw} = R_{sustained}/R_{peak}$ and L_{lat} to isolate infrastructure efficiency from algorithmic improvements.

- Significance:** The same ResNet-50 model can deliver very different throughput across hardware stacks, precision formats, batch sizes, and compiler configurations, yet still report the same ImageNet Top-1 accuracy. System benchmarks capture this implementation gap, which is invisible to algorithmic benchmarks that only report accuracy.
- Distinction:** Unlike algorithmic benchmarks (which vary model architectures and training procedures to improve convergence accuracy), system benchmarks hold the algorithm fixed and vary the implementation (kernel libraries, quantization formats, batch sizes, and hardware generations) to measure how efficiently the hardware-software stack executes the iron law’s $O/(R_{peak} \cdot \eta_{hw})$ term.
- Common pitfall:** A frequent misconception is that a system benchmark result generalizes across workloads. An accelerator that achieves high utilization on ResNet-50 (a compute-friendly vision workload) may achieve much lower utilization on a recommendation system (a memory-bandwidth-bound workload). System benchmarks are workload-specific; no single metric characterizes a hardware platform.

Effective benchmark interpretation requires knowing the performance characteristics of target hardware. Whether a specific AI workload is compute bound or memory-bound provides essential insight for optimization decisions. Computational intensity, measured as FLOP/byte,⁹ determines performance limits. Consider an NVIDIA A100 GPU with 312 TFLOP/s of FP16 Tensor Core performance (FP32 is 19.5 TFLOP/s) and 2.04 TB/s memory bandwidth (SXM variant). Dividing peak compute by peak bandwidth yields an arithmetic intensity threshold of 153 FLOP/byte.

7 | TPU (tensor processing unit): Google’s custom ASIC for neural network workloads (architecture details in chapter 11). A TPU v4 pod (4,096 chips) delivers 1.1 exaFLOP/s peak BF16 (Jouppi et al. 2023), but benchmarking TPUs requires caution: their systolic-array architecture favors regular tensor operations, so peak FLOP/s overstate performance on irregular workloads like sparse attention or dynamic control flow.

8 | ASIC (application-specific integrated circuit): An ASIC’s peak TOPS number applies only to the specific operators it was designed for. A single unsupported layer forces fallback to a general-purpose processor, potentially negating the entire efficiency advantage. This makes operator coverage the first question in any ASIC benchmark: the gap between peak and achieved throughput is not a hardware limitation but a workload-compatibility limitation.



Larger batches push transformer inference from memory-bound to compute-bound.

9 | FLOP/s (floating-point operations per second): The gap between advertised peak FLOP/s and achieved FLOP/s is the central tension in hardware benchmarking. The A100 advertises 312 TFLOP/s FP16 Tensor Core, but real workloads achieve different fractions of peak depending on arithmetic intensity, memory access patterns, precision, and runtime overhead. Reporting peak FLOP/s without utilization context is the most common benchmarking distortion.

Workloads below this threshold are bottlenecked by memory bandwidth, while those above are bottlenecked by compute capacity. The Roofline Model in section 11.6 provides the architectural foundation for interpreting these benchmark results. Section D.2.1 derives the roofline equation and the ridge-point threshold from first principles, so the arithmetic intensity bound used here can be reconstructed for any accelerator.

Roofline position¹⁰ depends on the workload. In this worked A100 example, an illustrative high-intensity ResNet-50 forward-pass workload at large batch size uses an arithmetic intensity of ~300 FLOP/byte, above the A100 ridge, and therefore represents compute-bound kernels (He et al. 2016a; Choquette et al. 2021). Lower-intensity operations fall below the ridge into the memory-bound regime: a BERT inference at batch size one, counting only weight-loading traffic, reaches ~100 FLOP/byte arithmetic intensity and a lower performance ceiling than peak. Increasing the batch size moves that same workload across the ridge from memory-bound to compute-bound (Pope et al. 2023). Section D.2.1.1 works the intensity-to-utilization calculation end to end on the A100, contrasting a compute-bound matrix multiplication kernel against a memory-bound element-wise kernel, so the steps generalize to any model-hardware pair.

A worked BERT inference estimate shows how these roofline principles translate into concrete deployment predictions.

¹⁰ | **Roofline model:** Williams et al. (2009) introduced the Berkeley model, named for the visual shape of its performance ceiling. Its ridge point (peak FLOP/s divided by peak bandwidth) separates memory-bound from compute-bound workloads, showing whether optimization should target data movement or arithmetic.

Napkin Math 12.2: Roofline analysis for BERT inference

Problem: BERT-Base must be deployed for inference on an A100 GPU. Management expects high GPU utilization. What performance should be predicted, and how can it be improved?

Step 1: Hardware limits.

- Peak compute: 312 TFLOP/s (FP16 Tensor Core)
- Memory bandwidth: 2.04 TB/s
- Ridge point: $312 \text{ TFLOP/s} \div 2.04 \text{ TB/s} = 153 \text{ FLOP/byte}$

Any workload with arithmetic intensity below 153 FLOP/byte is memory bound; above is compute bound.

Step 2: BERT-base characteristics.

- Parameters: 110M = 220 MB (FP16)
- FLOPs per inference: ~22 GFLOP (forward pass with sequence length $S = 128$)
- Data movement: ~220 MB (must load all weights from memory)
- Arithmetic intensity: $(22 \times 10^9) \div (220 \times 10^6) = 100 \text{ FLOP/byte}$ (weights-only model; see note in main text)

Step 3: Performance prediction. Since $100 \text{ FLOP/byte} < 153 \text{ FLOP/byte}$, BERT at batch = 1 is memory bound:

$$\text{Achievable perf} = 100 \text{ FLOP/byte} \times 2.04 \text{ TB/s} = 203.9 \text{ TFLOP/s}$$

$$\text{GPU utilization} = 203.9 \text{ TFLOP/s} / 312 \text{ TFLOP/s} = 65.4\%$$

Step 4: Optimization via batching. Increase batch size to 32:

- Same 220 MB of weights, but 32× more compute
- New FLOPs: $22 \times 10^9 \times 32 = 704 \text{ GFLOP}$
- New intensity: $(704 \times 10^9) \div (220 \times 10^6) = 3200 \text{ FLOP/byte}$

Since $3200 \text{ FLOP/byte} > 153 \text{ FLOP/byte}$, batch = 32 is compute bound. Assuming the implementation then sustains 85 percent of peak:

$$\text{Achievable perf} \approx 85\% \times 312 \text{ TFLOP/s} = 265.2 \text{ TFLOP/s}$$

$$\text{GPU utilization} \approx 85\%$$

Systems insight: Batch size can transform memory-bound inference into compute-bound inference by raising arithmetic intensity. The displayed 85 percent is an implementation-efficiency assumption, not a roofline prediction. Batching also increases latency because the system must wait to accumulate requests. This is the fundamental throughput-latency trade-off that MLPerf scenarios capture: SingleStream (batch = 1, latency-optimized) vs. Offline (maximum batch, throughput-optimized).

System benchmarks evaluate performance across scales, ranging from single-chip configurations to large distributed systems and covering both training and inference. Figure 12.1 juxtaposes sourced ImageNet top-5 error rates (Russakovsky et al. 2015; Krizhevsky et al. 2012) with a reconstruction of the growing use of GPUs in challenge entries, read from the entry counts NVIDIA charted for the challenge (Gray 2015). The two series show contemporaneous trends, not a causal estimate of how much of the accuracy improvement came from hardware rather than algorithms, data, or training practice.




Figure 12.1: ImageNet Error and GPU Adoption: Sourced top-5 error rates are shown beside GPU use among challenge entries. The temporal association does not isolate hardware’s causal contribution. Error-rate sources: (Russakovsky et al. 2015; Krizhevsky et al. 2012). Gray (2015) charts the per-year GPU entry counts.

The ImageNet example places GPU adoption alongside falling error rates without isolating either contribution; section 12.11.1 revisits this progression through model-specific architectural milestones. Effective system benchmarking, however, requires understanding the relationship between workload characteristics and hardware utilization. Modern AI systems rarely achieve theoretical peak performance due to interactions between computational patterns, memory hierarchies, and system architectures. This gap between theoretical and achieved performance shapes the design of meaningful system benchmarks.

Realistic hardware utilization patterns are essential for actionable benchmark design. As the preceding roofline analysis illustrated, GPU utilization varies with batch size and model architecture; the compute-bound value of 85 percent is an assumption, while the memory-bound single-request value is 65.4 percent. These patterns extend to memory bandwidth: parameter-heavy transformer inference and activation-heavy convolutional workloads stress different parts of the memory hierarchy, directly impacting achievable performance across different precision levels.

Effective system benchmarks must measure realistic utilization rather than peak theoretical capability, and this requirement establishes several scope boundaries. Energy is one dimension: performance per watt varies widely across platforms, and an underutilized accelerator consumes disproportionate power for its output, penalizing both operational cost and environmental impact. Distribution is another: multi-node training adds communication bottlenecks, network-topology effects, and coordination overhead that single-node benchmarks cannot capture and that warrant dedicated treatment beyond this book. Within the single-machine scope here, multi-GPU benchmarking instead focuses on intra-node communication, memory-bandwidth utilization across accelerators, and gradient-synchronization efficiency across distinct accelerator memories connected by NVLink

or PCIe. Across all of these, a benchmark earns its value only when its operating point matches the deployment's, not the datasheet's.

12.3.3 Community-driven standardization

The hardware utilization insights are only useful for comparison when measured consistently, which requires community-driven standardization. When one team measures inference latency with preprocessing included and another excludes it, when accuracy benchmarks use different data splits, or when power measurements employ different system boundaries, meaningful comparison becomes impossible. Individual organizations cannot establish measurement standards alone; the proliferation of benchmarks across the three dimensions creates fragmentation that only coordinated effort can resolve.

The most successful benchmarks emerge through broad collaboration among academic institutions, industry partners, and domain experts. ImageNet's lasting impact demonstrates how sustained community engagement through workshops, challenges, and open datasets establishes authority that corporate-driven benchmarks rarely achieve. This collaborative development creates a foundation for formal standardization: IEEE working groups (IEEE Standards Association 2024) and ISO/IEC technical committees (ISO 2024) codify community-developed methodologies into official standards (for example, IEEE 2416 (IEEE Standards Association 2019a) for system power modeling), providing precise measurement specifications that enable reliable cross-institutional comparison. Projects that provide open-source reference implementations, containerized evaluation environments, and comprehensive validation suites further reduce barriers and ensure consistent interpretation across research groups.

ML benchmarks must balance academic rigor with industry practicality, since theoretical advances must translate to practical improvements in deployed systems (Mattson et al. 2020; Reddi et al. 2019). Benchmarks that emerge from this balance, with transparent governance and regular evolution, become durable reference points; those developed in isolation struggle to gain traction regardless of technical sophistication. These evaluation methodology principles guide both training and inference benchmark design throughout this chapter.

Community standards ensure reproducibility, but they do not prescribe the level of detail at which measurements should be taken. A benchmark could time a single matrix multiplication or an entire training run—and each choice reveals different kinds of information. The depth of measurement, from individual operations to complete systems, determines what insights benchmarks can provide and which problems they can diagnose.

12.4 Benchmarking Granularity

A GPU kernel that runs $3\times$ faster in isolation may deliver zero end-to-end speedup if the data pipeline cannot keep pace. This diagnostic failure makes evaluation granularity a fundamental design choice. Standardization specifies *how* measurement is consistent, while **benchmarking granularity** specifies *what* is measured. Each validation dimension can be assessed at different scales, from individual operations to complete workflows, with each granularity level revealing different kinds of problems:

- Micro benchmarks isolate individual components: kernel execution time, memory bandwidth utilization, single-layer accuracy. These diagnose *where* problems occur.
- Macro benchmarks evaluate subsystems: full model training convergence, inference pipeline throughput, dataset bias metrics. These reveal *what* problems exist.
- End-to-end benchmarks measure complete workflows: request-to-response latency including preprocessing, training time-to-accuracy including data loading, model performance on production data distributions. These show whether the system works.

The optimization techniques from Part III operate at different granularities (kernel fusion targets micro performance, pruning affects macro model behavior, data curation determines end-to-end generalization) and validation must match. A micro benchmark might show kernel speedup while a macro benchmark reveals memory bottlenecks that negate the gain; an end-to-end benchmark might expose data pipeline stalls invisible at any other level.

Figure 12.2 maps these granularity levels onto the ML stack by breaking the stack into four distinct evaluation scopes. Each scope progressively expands the measurement boundary. Micro-benchmarks isolate neural network layers, macro-benchmarks encompass complete models, application benchmarks add supporting compute, and end-to-end benchmarks capture the full deployment context including non-AI components.




Figure 12.2: Benchmarking Granularity: Four-panel block diagram showing micro, macro, application, and end-to-end evaluation layers. Each panel maps a distinct scope of assessment, from isolated kernel operations through full-system deployment, enabling targeted optimization at every level of the ML stack.

12.4.1 Micro benchmarks

While end-to-end benchmarks reveal overall system behavior, optimization requires pinpointing exactly which operations consume time and energy. Micro-benchmarks serve this diagnostic purpose by isolating individual tensor operations, the mathematical primitives optimized in chapter 11.

Consider debugging a slow inference pipeline: macro benchmarks might show unacceptable latency, but only micro-benchmarks reveal whether the bottleneck lies in convolutions, attention mechanisms, memory copies, or activation functions. That isolation connects theoretical hardware capabilities to realized performance by showing which operation owns the cost.

🔍 Systems Perspective 12.4: Micro-benchmarking rules

To avoid measuring hardware artifacts instead of kernel performance, follow the *Systems Detective's Rules*:

1. **The warm-up rule:** Do not measure cold-start iterations as steady-state performance. Modern hardware uses DVFS (dynamic voltage and frequency scaling) and dynamic frequency boosting; caches, kernels, and clocks need warm-up before the measured loop represents sustained behavior.
2. **The variance rule:** Report the coefficient of variation (CV) ($CV = \sigma_{\text{run}} / \mu_{\text{run}}$), where σ_{run} and μ_{run} are the standard deviation and mean across repeated runs. Investigate values above the protocol's workload-specific tolerance; common causes include background OS jitter, thermal throttling, and memory contention, and 5 percent is not a universal cutoff.
3. **The “speed of light” (SOL) check:** Compare the achieved throughput against the roofline. If a kernel achieves 10 TFLOP/s on an H100 (peak ~989 TFLOP/s FP16, or ~1,979 TFLOP/s FP8 dense), the diagnostic step is to identify the cause of low utilization (often kernel launch latency from too many small kernels) before optimizing the code itself.
4. **The cache-state rule:** Set cache state to match the claim. Flush or exceed the L2 cache when measuring cold dynamic random-access memory (DRAM) bandwidth; preserve a warmed cache when measuring steady-state reuse. Otherwise a nominal DRAM result may instead reflect cache bandwidth (~5 TB/s–10 TB/s) rather than DRAM bandwidth (~1 TB/s–2 TB/s).

A key area of micro-benchmarking focuses on tensor operations, the computational core of deep learning. Libraries like cuDNN¹¹ (Chetlur et al. 2014) by NVIDIA provide optimized primitives for core computations such as convolutions and matrix multiplications across different hardware configurations. Micro-benchmarks around these primitives help developers understand how their hardware handles the core mathematical operations that dominate ML workloads.

¹¹ | **cuDNN (CUDA deep neural network library):** Released by NVIDIA in 2014, cuDNN provides hand-tuned kernel implementations for convolutions, pooling, and normalization. The benchmarking implication: reported inference latencies depend heavily on which cuDNN version and algorithm autotuner settings were used, making cuDNN version a mandatory element of any reproducible benchmark specification.

Measuring these operations correctly requires discipline. A small set of measurement rules prevents common errors that can invalidate results entirely.

A profiler turns these measurement rules into iron-law evidence by decomposing execution time into data movement, compute throughput, and latency overhead, as introduced in section 1.7.

Systems Perspective 12.5: Measuring the iron law terms

Moving from theory to trace means mapping the iron law equation from section 1.7 onto a profiler timeline (like Nsight Systems or PyTorch Profiler).

Measuring the data term ($\frac{D_{\text{vol}}}{\text{BW}}$)

- **Signal:** Look for the “Memory Throughput” or “DRAM Bandwidth” line.
- **Formula:** $\text{BW}_{\text{effective}} = \frac{D_{\text{vol}}}{T_{\text{kernel}}}$.
- **Diagnosis:** If $\text{BW}_{\text{effective}} \approx \text{BW}_{\text{peak}}$ (for example, close to the A100’s 2.04 TB/s peak), the kernel is memory bound. Optimizing compute (O) will do nothing.

Measuring achieved compute throughput ($R_{\text{peak}} \cdot \eta_{\text{hw}}$)

- **Signal:** Look for “SM Active” or “Compute Throughput”.
- **Formula:** Achieved TFLOP/s = $\frac{O}{10^{12} T_{\text{kernel}}}$.
- **Diagnosis:** If Achieved TFLOP/s $\ll$ Peak TFLOP/s AND $\text{BW}_{\text{effective}} \ll \text{BW}_{\text{peak}}$, the system is in the “Utilization Trap”: likely Latency Bound (kernels too small) or Grid Bound (not enough threads).

Measuring the latency term (L_{lat})

- **Signal:** Look for gaps (empty space) between colored kernel bars on the timeline.
- **Formula:** Overhead Ratio = $\frac{T_{\text{gap}}}{T_{\text{kernel}} + T_{\text{gap}}}$.
- **Diagnosis:** A “Sawtooth” pattern (Compute, Gap, Compute, Gap) indicates high software overhead. The solution is operator fusion, covered in section 11.8.1.3, or CUDA Graphs, which capture a repeated sequence of GPU launches so the runtime can replay it with less CPU dispatch overhead.

While benchmarks like MLPerf reveal *how fast* a system is, micro-benchmarking tools reveal *why* it is slow. To perform this diagnosis, engineers use kernel-level profilers that peer inside the execution of individual operations.

Framework profilers

Tools like PyTorch Profiler capture the logical execution flow of a training or inference step. They identify which layer dominates runtime, whether CPU and GPU work overlap or synchronize unnecessarily, and whether the data loader keeps the accelerator supplied. The diagnostic metric is the step-time breakdown across data loading, compute, and communication, because that breakdown tells the engineer which subsystem owns the next optimization.

Kernel profilers

Tools like NVIDIA Nsight Systems and Compute capture physical execution on the hardware. They determine whether a matrix multiplication is compute bound or memory bound, whether the Streaming Multiprocessors reach high occupancy, and whether memory accesses obey coalescing rules. The diagnostic metric is roofline position, because FLOP/s relative to memory bandwidth reveals whether more arithmetic throughput can help or whether the kernel is waiting on data movement.

The recommended workflow is to start with the Framework Profiler to find the slow layer (for example, “The Attention Block is slow”). Then, use the Kernel Profiler to diagnose the physics (for example, “The Softmax kernel is memory bound because it is reading too many bytes per FLOP”). This targeted approach avoids the “optimization without measurement” trap.

Micro-benchmarks also examine activation functions and neural network layers in isolation. This includes measuring the performance of various activation functions like the rectified linear unit (ReLU), Sigmoid, and Tanh under controlled conditions, and evaluating the computational efficiency of distinct neural network components such as long short-term memory cells or transformer blocks when processing standardized inputs.

DeepBench (Baidu Research 2016), developed by Baidu, was one of the first to demonstrate the value of comprehensive micro-benchmarking. It evaluates these core operations across different hardware platforms, providing detailed performance data that helps developers optimize their deep learning implementations. By isolating and measuring individual operations, DeepBench enables precise comparison of hardware platforms and identification of potential performance bottlenecks.

These granular measurements enable precise optimization, but they cannot reveal how components interact when assembled into complete models. Macro-benchmarks address this gap.

12.4.2 Macro benchmarks

Micro-benchmarks confirm that individual convolution kernels run fast. Macro-benchmarks reveal whether the complete model works under realistic conditions. This shift from component-level to model-level assessment reveals how architectural choices and component interactions affect overall model behavior. For instance, while micro-benchmarks might show optimal performance for individual convolutional layers, macro-benchmarks reveal how these layers work together within a complete convolutional neural network.

Macro-benchmarks exist to serve one decision: choosing a model or architecture under standardized conditions. That decision needs the performance dimensions that emerge only at the model level: prediction accuracy, which shows how well the model generalizes to new data; memory consumption patterns across different batch sizes and sequence lengths; throughput under varying computational loads; and latency across different hardware configurations. These dimensions interact in ways a single-layer micro-benchmark cannot expose. A model that wins on accuracy may lose once its memory footprint at the target sequence length forces a smaller batch, collapsing the throughput that made it attractive, a coupling visible only when the complete model is measured as a unit.

The assessment of complete models occurs under standardized conditions using established datasets and tasks. For example, computer vision models might be evaluated on ImageNet (Deng et al. 2024), measuring both computational efficiency and prediction accuracy. Natural language processing models might be assessed on translation tasks, examining how they balance quality and speed across different language pairs.

Several industry-standard benchmarks make model-level comparison reproducible across platforms. The MLPerf family (Inference, Mobile, Client, and Tiny) provides comprehensive testing suites adapted for computational environments from data center to microcontroller, detailed in section 12.8.4. For embedded systems, EEMBC's MLMark emphasizes both performance and power efficiency, while the AI-Benchmark (Ignatov and Timofte 2024) suite specializes in mobile platforms.

12.4.3 End-to-end benchmarks

End-to-end benchmarks provide the most inclusive evaluation by encompassing the entire pipeline of an AI system, not just the model. This includes extract, transform, load data processing; model inference; postprocessing of results; and critical infrastructure components like storage and network systems.

Data processing (extracting from source systems, transforming through cleaning and feature engineering, and loading into model-ready formats) forms the foundation of the pipeline. These preprocessing steps directly affect overall performance, and end-to-end benchmarks must assess standardized datasets through complete pipelines to ensure data preparation does not become a bottleneck. Postprocessing similarly affects real-world performance: a computer vision system must postprocess detection boundaries, apply confidence thresholds, and format results for downstream applications before the user sees a response.

Infrastructure components heavily influence overall performance beyond the AI workload itself. Storage solutions can dominate data retrieval times with large AI datasets, and network interactions in distributed systems can become performance bottlenecks. End-to-end benchmarks must evaluate

these components under specified environmental conditions to ensure reproducible measurements of the entire system.

Public end-to-end benchmarks rarely account for data storage, network, and compute performance in one measurement. While MLPerf Training and Inference approach end-to-end evaluation, they primarily focus on model performance rather than real-world deployment scenarios. Nonetheless, they provide valuable baseline metrics for assessing AI system capabilities.

Given the inherent specificity of end-to-end benchmarking, organizations typically perform these evaluations internally by instrumenting production deployments. The sensitivity of these measurements means they rarely appear publicly, but their absence from the literature does not diminish their importance.

12.4.4 Granularity trade-offs and selection criteria

Table 12.4 reveals how different challenges emerge at different stages of an AI system's lifecycle. Each benchmarking approach provides unique insights: micro-benchmarks help engineers optimize specific components like GPU kernel implementations or data loading operations, macro-benchmarks guide model architecture decisions and algorithm selection, while end-to-end benchmarks reveal system-level bottlenecks in production environments.

Table 12.4: Benchmarking Granularity Levels: Different benchmark scopes target distinct stages of ML system development. Micro-benchmarks isolate individual operations for low-level optimization, macro-benchmarks evaluate complete models to guide architectural choices, and end-to-end benchmarks assess full system performance in production environments.

Component	Micro Benchmarks	Macro Benchmarks	End-to-End Benchmarks
Focus	Individual operations	Complete models	Full system pipeline
Scope	Tensor ops, layers, activations	Model architecture, training, inference	Extract, transform, load; model; infrastructure
Example	Conv layer performance on cuDNN	ResNet-50 on ImageNet	Production recommendation system
Advantages	Precise bottleneck identification, Component optimization	Model architecture comparison, Standardized evaluation	Realistic performance assessment, System-wide insights
Challenges	May miss interaction effects	Limited infrastructure insights	Complex to standardize, Often proprietary
Typical Use	Hardware selection, Operation optimization	Model selection, Research comparison	Production system evaluation

Picking a single granularity level is rarely sufficient because a core tension exists between diagnostic precision and real-world fidelity. Figure 12.3 maps this trade-off, placing micro-benchmarks at the high-isolation end (precise but narrow) and end-to-end benchmarks at the high-representativeness end (realistic but harder to diagnose). No single point on this spectrum provides both: micro-benchmarks pinpoint exactly which kernel is slow but miss system-level bottlenecks, while end-to-end benchmarks capture production behavior but obscure root causes. The practical takeaway is that effective ML system evaluation requires combining insights from all three levels.

Component interaction often produces unexpected behaviors that single-level benchmarks miss. While micro-benchmarks might show excellent performance for individual operations and macro-benchmarks might demonstrate strong model accuracy, end-to-end evaluation can reveal that data preprocessing creates unexpected bottlenecks during high-traffic periods. These system-level insights remain hidden when components undergo isolated testing.

Choosing a granularity level, however, is only half the design problem. The other half is specifying the concrete ingredients every benchmark requires: the task, data, model, and metrics. Without those ingredients, even the right granularity level produces meaningless numbers. The components of a benchmark determine whether results translate into actionable engineering insight or merely generate impressive-looking numbers that collapse under scrutiny.

12.5 Benchmark Components

Choosing between micro, macro, and end-to-end granularity determines what a benchmark can diagnose, but every benchmark at every granularity must still specify the task, data, model, metrics, harness, system context, and run rules that make its result interpretable. Micro-benchmarks




Figure 12.3: Isolation vs. Representativeness: The core trade-off in benchmarking granularity. Micro-benchmarks provide high diagnostic precision but limited real-world relevance, while end-to-end benchmarks capture realistic system behavior but offer less precise component-level insights. Effective ML system evaluation requires strategic combination of all three levels.

require synthetic inputs that isolate specific computational patterns; macro-benchmarks demand representative datasets like ImageNet; end-to-end benchmarks must incorporate real-world data with all its noise and distributional shift. Despite this variation, all benchmarks share a common implementation problem: each component must constrain the next one so the final number has a defensible meaning.

The essential components interconnect to form a complete evaluation pipeline. The workflow in figure 12.4 traces nine stages of an industrial audio anomaly detection benchmark, from problem definition through quantization to ARM embedded deployment. The serial dependency is the critical observation: the task definition constrains which datasets are valid, the dataset properties determine which model architectures are feasible, and the target hardware dictates quantization and compilation choices. Anomaly detection serves as an effective illustration precisely because it spans the full stack, coupling ML inference accuracy with embedded systems constraints such as memory footprint, power budget, and real-time latency. A benchmark that measured only classification accuracy or only inference speed would miss the interactions between these stages, where a decision at any point propagates forward and narrows every subsequent choice.




Figure 12.4: Benchmark Workflow Components: Read left to right: each stage fixes assumptions for the next, carrying an audio anomaly-detection task from problem definition to a reproducible embedded-device benchmark run.

Effective benchmark design must account for the optimization techniques established in preceding chapters. Quantization and pruning affect model accuracy-efficiency trade-offs, requiring benchmarks that measure both speedup and accuracy preservation simultaneously. Hardware acceleration

techniques influence arithmetic intensity and memory bandwidth utilization, necessitating Roofline Model analysis to interpret results correctly. Understanding these optimization foundations enables benchmark selection that validates claimed improvements rather than measuring artificial scenarios.

12.5.1 Problem definition

Every benchmark begins by specifying exactly what the system must do. The anomaly detection system in figure 12.4 processes audio signals to identify deviations from normal operation patterns, an industrial monitoring application that exemplifies how formal task specifications translate into practical implementations. While specific tasks vary widely by domain (natural language processing tasks include machine translation, question answering (Hirschberg and Manning 2015), and text classification; computer vision employs object detection and image segmentation (Everingham et al. 2009; Lin et al. 2014)), every benchmark task specification must define three essential elements: an input specification (what data the system processes), an output specification (what response the system must produce), and a performance specification (quantitative requirements for accuracy, speed, and resource utilization).

Task design directly impacts the benchmark’s ability to evaluate AI systems. The audio anomaly detection example illustrates this through its specific requirements: processing continuous signal data, adapting to varying noise conditions, and operating within strict time constraints. These practical constraints create a framework for assessment that reflects real-world operational demands. Each subsequent phase of benchmark implementation, from dataset selection through deployment, builds directly upon these initial specifications.

12.5.2 Standardized datasets

A task definition is only as good as the data used to evaluate it. Standardized datasets ensure that all models undergo testing under identical conditions, enabling direct comparisons across different approaches—without them, every team would evaluate on private data, making cross-lab comparison impossible. In computer vision, ImageNet (Deng et al. 2024, 2009), COCO (Lin et al. 2014), and CIFAR-10 (Krizhevsky 2009) serve as reference standards; in natural language processing, SQuAD¹² (Rajpurkar et al. 2016), GLUE¹³ (A. Wang et al. 2018), and WikiText (Merity 2016; Merity et al. 2016) fulfill similar roles, each encompassing a range of complexities and edge cases.

Dataset selection is the first place a benchmark can lose contact with deployment reality. In the audio anomaly detection example (figure 12.4), the dataset must include representative waveform samples of normal operation alongside comprehensive examples of anomalous conditions. Domain-specific collections cover different audio tasks: ToyADMOS¹⁴ (Koizumi et al. 2019) supports controlled anomaly-detection research, while Google Speech Commands (Warden 2018) supports keyword recognition. Effective benchmark datasets must balance two competing demands: accurately representing real-world challenges while maintaining sufficient complexity to differentiate model performance. Simplified datasets like ToyADMOS are valuable for methodological development but may not capture the full complexity of production environments.

12.5.3 Model selection

With task and data specified, the benchmark must define which models to evaluate and what baselines to compare against. This choice is less straightforward than it appears: a benchmark’s model selection determines whether results reflect architectural innovation, implementation quality, or simply framework-specific optimizations. The selection process builds upon the architectural foundations established in chapter 6 and must account for the framework considerations discussed in chapter 7.

Baseline models serve as reference points spanning from basic implementations (linear regression, logistic regression) to advanced architectures with proven success in comparable domains. In NLP, models like BERT¹⁵ have emerged as standard baselines. Critically, the choice of baseline depends on the deployment framework: a PyTorch implementation may exhibit different performance characteristics than its TensorFlow equivalent due to framework-specific optimizations and operator implementations, meaning the benchmark must control for this variable.

Once the architecture is selected, model development follows two parallel optimization paths that the benchmark must track. Training optimization focuses on achieving target accuracy within computational constraints. Inference optimization addresses the transition to production—particularly

¹² | **SQuAD (Stanford question answering dataset):** Introduced in 2016 with more than 100,000 question-answer pairs from Wikipedia (Rajpurkar et al. 2016). AI systems exceeded the SQuAD 1.1 human baseline of 91.2 percent F1 by 2018 (Devlin et al. 2019), but this “superhuman” result illustrates a benchmarking failure mode: the task’s extractive format (answers are text spans within the passage) makes it easier than open-ended question answering, inflating perceived capability relative to production NLP systems.

¹³ | **GLUE (general language understanding evaluation):** Introduced in 2018 as a broad language-understanding benchmark (A. Wang et al. 2018), GLUE was quickly saturated by systems such as BERT (Devlin et al. 2019). This is Goodhart’s Law in action: once GLUE became a target, leaderboard optimization reduced its discriminating power. The pattern motivated harder follow-on evaluations such as SuperGLUE and BIG-bench.

¹⁴ | **ToyADMOS:** Developed by NTT Communications in 2019 for acoustic anomaly detection, containing audio recordings from toy car, toy conveyor, and related miniature-machine operating sounds (Koizumi et al. 2019). The “toy” prefix is intentional: the controlled environment enables reproducible benchmarking but can create a domain gap when models are moved to noisier industrial environments with different machines, sensors, vibration, and background sound.

¹⁵ | **BERT (bidirectional encoder representations from transformers):** BERT-Large (340M parameters) is a language-processing workload in MLPerf Inference. The benchmark fixes the task, data, quality target, and applicable execution scenarios so systems can be compared on the same workload. BERT latency still depends on sequence length, batching, and implementation.

precision reduction from FP32 to INT8 or lower, which demands careful calibration to maintain accuracy while reducing resource requirements. The benchmark must specify requirements for both paths, because a model that trains efficiently but deploys poorly (or vice versa) fails the full evaluation. This dual optimization naturally demands quantitative evaluation metrics that span all three dimensions of the benchmarking framework.

12.5.4 Evaluation metrics

Evaluation metrics¹⁶ translate raw model behavior into numbers that can be compared, ranked, and used to make engineering decisions. The challenge is choosing the right numbers: a metric that captures accuracy but ignores latency may declare the winner to be a model too slow for production; one that rewards throughput but ignores energy may optimize for a deployment budget that does not exist.

Table 12.5 should be read as a decision aid: it categorizes metrics by the failure mode each exposes and the deployment context it serves.

Table 12.5: ML Benchmarking Metric Taxonomy: Metrics organized by evaluation category, unit, and primary use case. Accuracy metrics quantify model quality, throughput and latency metrics capture system speed, and efficiency metrics combine multiple dimensions. Selecting the right metric for the deployment context is often more important than optimizing any single metric to its maximum.

Category	Metric	Unit	Primary Use Case
Accuracy	Top-1/Top-5 Accuracy	Percentage	Classification
	mAP (mean Average Precision)	0–1 score	Object detection
	BLEU/ROUGE	0–100 score	NLP generation
	Perplexity	Score (lower = better)	Language modeling
Throughput	Samples/second	Samples/s	Batch inference
	Token throughput	tokens/s	LLM inference
	Time-to-train	Hours/days	Training benchmarks
Latency	p50 latency	Milliseconds	Median response time
	p99 latency	Milliseconds	Tail latency (SLA)
	First-token latency	Milliseconds	LLM responsiveness
Efficiency	Samples/second/watt	Samples/s/W	Energy efficiency
	Accuracy/FLOP	percent/PFLOP	Algorithmic efficiency
	TCO per inference	\$/inference	Economic efficiency

Several distinctions within this taxonomy deserve emphasis. Throughput measures aggregate capacity (ideal for batch processing), while latency measures individual request timing (critical for interactive applications). These metrics frequently conflict: maximizing throughput through batching often increases per-request latency. Mean latency can hide problematic tail behavior—a system with 10 ms mean latency might have 500 ms p99 latency, failing SLA requirements. In production, percentiles (p50, p95, p99) are far more informative than means. Finally, compound metrics like samples/second/watt combine multiple dimensions into a single number, enabling quick comparisons but obscuring individual bottlenecks. Reporting both atomic and compound metrics provides a complete picture.

Metric choice must align with task objectives and deployment constraints, because the same raw model behavior can produce different scores across frameworks. The training methodologies from chapter 8 demonstrate how different frameworks handle loss computation and gradient accumulation differently, affecting reported metrics. Even small implementation differences, such as evaluation-mode batch-normalization handling, can shift measured accuracy enough to matter when benchmark deltas are small.

Task-specific metrics quantify a model’s performance on its intended function. For example, classification tasks employ metrics including accuracy (overall correct predictions), precision (positive prediction accuracy), recall (positive case detection rate), and F1 score (precision-recall harmonic mean) (Sokolova and Lapalme 2009). Regression problems use error measurements like Mean Squared Error (MSE) and Mean Absolute Error (MAE) to assess prediction accuracy. Domain-specific applications often require specialized metrics; for example, machine translation uses BLEU¹⁷

¹⁶ | **Metric:** In mathematics, a metric is a distance function satisfying strict axioms including the triangle inequality. ML borrows the term loosely for quantitative measures such as BLEU and perplexity, which are scoring rules rather than mathematical metrics. Leaderboard rankings can change when the evaluation protocol, dataset slice, or metric weighting changes, making the choice of metric an engineering decision that shapes which system wins, not just how we measure it.

¹⁷ | **BLEU (bilingual evaluation understudy):** Introduced by IBM in 2002, BLEU measures translation quality through modified n-gram precision with a brevity penalty against reference translations (Papineni et al. 2002). BLEU is a canonical example of Goodhart’s Law in ML: optimizing for n-gram matches can reward surface-level word overlap even when meaning, fluency, or deployment usefulness diverges from the target.

to measure modified n-gram precision against one or more human reference translations (Papineni et al. 2002).

Production deployment adds implementation metrics to task metrics. Model size, measured in parameters or memory footprint, directly affects deployment feasibility across different hardware platforms. Processing latency, typically measured in milliseconds per inference, determines whether the model meets real-time requirements. Energy consumption, measured in watts or joules per inference, indicates operational efficiency. These practical considerations reflect the growing need for solutions that balance accuracy with computational efficiency. The operational challenges of maintaining these metrics in production environments are explored in deployment strategies (chapter 14).

The benchmark therefore needs a metric set that matches both task requirements and deployment constraints. A single metric rarely captures all relevant aspects of performance in real-world scenarios. For instance, in anomaly detection systems, high accuracy alone may not indicate good performance if the model generates frequent false alarms. Similarly, a fast model with poor accuracy fails to provide practical value.

This multi-metric evaluation approach appears in the anomaly detection system, which reports performance across multiple dimensions: model size (270K parameters), processing speed (10.4 ms/inference), detection accuracy (0.86 AUC), and energy consumption (516 μ J per inference). This combination of metrics ensures the model meets both technical and operational requirements in real-world deployment scenarios.

12.5.5 Benchmark harness

Metrics define *what* to measure; the benchmark harness determines *how* to measure it. A harness is the test infrastructure that delivers inputs to the system under test, collects measurements, and ensures that the entire process is reproducible. Without a well-designed harness, even perfectly chosen metrics produce unreliable numbers.

Harness design must align with the intended deployment scenario. For server deployments, the harness generates request patterns that simulate real-world traffic, often using a Poisson distribution¹⁸ to model random but statistically consistent workloads, while managing concurrent requests and varying load intensities.

For embedded and mobile applications, the harness generates input patterns that reflect actual deployment conditions. This might involve sequential image injection for mobile vision applications or synchronized multi-sensor streams for autonomous systems. Such precise input generation and timing control ensures the system experiences realistic operational patterns, revealing performance characteristics that would emerge in actual device deployment.

The harness must also accommodate different throughput models. Batch processing scenarios require the ability to evaluate system performance on large volumes of parallel inputs, while real-time applications need precise timing control for sequential processing. In the embedded implementation phase, the harness must support precise measurement of inference time and energy consumption per operation.

Reproducibility demands that the harness maintain consistent testing conditions across different evaluation runs. This includes controlling environmental factors such as background processes, thermal conditions, and power states that might affect performance measurements. The harness must also provide mechanisms for collecting and logging performance metrics without measurably impacting the system under test.

12.5.6 System specifications

Complementing the harness that controls test execution, system specifications document the complete computational environment: the hardware and software stack on which the benchmark runs. Without precise specifications, a reported throughput number is meaningless: the same model can train much faster on a newer accelerator than on an older one, making the hardware context inseparable from the result.

On the hardware side, specifications must capture the processor type and clock rate, accelerator model and memory (GPU, TPU, or custom ASIC), system RAM, storage type, and network configuration for distributed setups. On the software side, they must record the operating system,

18 | Poisson distribution: Named after Siméon Denis Poisson, who formalized it in 1837 while modeling wrongful conviction rates in French courts. The distribution models independent events at a constant average rate (λ_{arr}) and is a common baseline for server request arrivals. Real ML serving traffic often violates its assumptions through burstiness, correlation, and time-varying rates. A Poisson harness can therefore misestimate tail latency unless traces or stress cases cover those production patterns.

framework versions (for example, PyTorch 2.1 vs. TensorFlow 2.14), compiler flags, and environment management tools such as Docker containers or virtual environments. This level of detail enables other researchers to replicate the benchmark environment with high fidelity and provides critical context for interpreting performance differences.

Many benchmarks include results across multiple hardware configurations, precisely because the trade-offs between model complexity, computational resources, and performance only become visible through comparative analysis. As the field increasingly prioritizes sustainability, specifications now extend to energy consumption metrics such as FLOP/s per watt and total power draw over training time, reflecting growing awareness that computational efficiency is an engineering requirement, not merely an environmental aspiration.

12.5.7 Run rules

System specifications describe *what* the benchmark runs on; run rules govern *how* it runs. These procedural constraints make results interpretable and repeatable, which is harder than it sounds in a field where stochastic processes (weight initialization, data shuffling, and dropout masks) mean that two runs on identical hardware can produce different numbers. A protocol may fix seeds and data order when isolating system effects, or require repeated runs and quality thresholds when stochastic variation is part of the workload. In either case, the seed policy and known sources of nondeterminism must be explicit.

Hyperparameter documentation is equally critical. A learning-rate change can shift convergence and final accuracy, so a reproducible result records every configuration setting that can affect the outcome. Dataset versions, splits, and preprocessing must likewise be identified. When privacy or licensing prevents sharing data directly, the report must state that limitation and preserve enough provenance and transformation detail to judge comparability.

Code provenance completes the reproducibility chain. Strong benchmark protocols preserve the implementation version—not just the model, but the relevant preprocessing, training, and evaluation code—and disclose whether that code can be shared. Reference suites may distribute containerized environments that encapsulate dependencies and configurations, while experimental logs retain training metrics, checkpoints, and any mid-run adjustments. Together, these records turn a one-time measurement into evidence another team can inspect and, when access permits, reproduce.

12.5.8 Result interpretation

Producing benchmark numbers is the easy part; interpreting them correctly is where most engineers go wrong. A raw throughput figure or accuracy score is meaningless without understanding the conditions that produced it, the statistical confidence behind it, and the deployment context that determines whether the number matters.



Example 12.2: Benchmarking a vision model for edge deployment

Scenario: An edge AI team evaluates MobileNetV2 for a wildlife camera trap on a Raspberry Pi 4 under strict latency and model size limits (Sandler et al. 2018).

Diagnosis: FP32 execution yields 120 ms latency (8.3 FPS) and 14 MB size. INT8 quantization speeds up inference by 3.4× to 35 ms (28.6 FPS) and reduces model size by 4×, at a cost of 0.9 percentage points Top-1 accuracy.

Systems lesson: Edge model benchmarking requires evaluating multi-dimensional trade-offs across latency, accuracy, and memory payload. Quantization enables high-throughput edge execution when modest accuracy drops meet product requirements.

Before drawing conclusions from benchmark results, apply the vendor claim analysis framework introduced earlier (see the “Decoding Vendor Benchmark Claims” checklist) and extend it with two additional checks. First, the comparison must be fair: comparing ResNet-50 against MobileNet conflates architecture differences with optimization choices; precision differences (FP32 vs. INT8) can materially affect performance, and batch size, hardware generation, and software framework must all be controlled. Second, the statistics must be meaningful: reliable results require multiple runs, reported variance with confidence intervals, clear handling of outliers, and steady-state operation

rather than cold-start effects. Applying these questions to a representative vendor claim illustrates how incomplete specifications obscure real performance.

Beyond vendor claims, context determines which metrics matter most. A 1 percent accuracy improvement may be decisive for medical diagnostics but irrelevant for an application that prioritizes inference speed. Practitioners should also guard against *benchmark overfitting*, where models are excessively optimized for specific benchmark tasks at the expense of real-world generalization, by evaluating performance on related but distinct tasks and considering practical deployment scenarios.

Systems Perspective 12.6: Interpreting a benchmark claim

A vendor claims “Our system achieves 10,000 inferences/second on ResNet-50.” Taken alone, the number cannot support deployment planning because the conditions that determine its meaning are unstated.

Four unstated conditions determine what the claim means:

1. Batch size: Large batches often achieve high throughput but can violate latency targets; batch 1 achieves low latency but lower throughput.
2. Precision: INT8 performance relative to FP32 depends on the model, operators, hardware, and implementation, and may have accuracy or calibration implications.
3. Measurement boundary: The claim must state whether it covers pure inference or includes preprocessing.
4. Accuracy: The claim must state whether the model matches the original 76.1 percent Top-1 or a degraded level.

Example: “10,000 inferences/second on ResNet-50 at batch size 32, INT8 precision, 76 percent Top-1 accuracy, including JPEG decoding, on NVIDIA H100 at 700 W thermal design power (TDP).”

Systems insight: Understanding whether a performance difference is meaningful requires both statistical rigor and contextual validation. A benchmark number without these details is a marketing claim, not an engineering specification.

12.5.9 Example benchmark

To see how these components work together in practice, walk through the anomaly detection pipeline in figure 12.4 one more time, now focusing on the output stage. The benchmark produces three complementary measurements: a model size of 270K parameters with 10.4 ms per inference (computational resources), a detection accuracy of 0.86 AUC in distinguishing normal from anomalous audio patterns (task effectiveness), and an energy consumption of 516 μ J per inference (operational efficiency).

Which of these metrics matters most depends entirely on the deployment context. Energy per inference is especially critical for battery-powered devices and still affects server operating cost and power capacity. Model size constrains embedded devices with limited memory and determines accelerator capacity and replication cost in cloud deployments. Processing speed determines whether the system can operate in real-time or must batch inputs. These metrics also reveal inherent trade-offs: reducing model size from 270K parameters might improve speed and energy efficiency but degrade the 0.86 AUC detection accuracy. Whether these measurements constitute a “passing” benchmark depends on the deployment constraints—the framework provides structure for consistent evaluation, but acceptance criteria must come from the application requirements.

The components just enumerated define how to assemble any single benchmark. Two benchmark categories recur often enough across the optimization pipeline to warrant their own component checklists here: compression benchmarks, which a pruned or quantized model must pass before deployment, and mobile and edge benchmarks, which a power- and thermally-constrained target imposes. Each composes the task, data, model, metrics, harness, and run rules just defined while adding constraints the generic checklist does not. Both are previews of dimensions the chapter develops fully later: compression validation returns in section 12.11.1.3 with the full multi-metric protocol, and sustained-power behavior returns in section 12.9.

12.5.10 Compression benchmarks

Neural network compression (pruning, quantization, knowledge distillation, and architecture optimization) requires specialized benchmarks because compression reshapes the trade-off landscape: every byte saved or operation eliminated must be weighed against potential accuracy loss and hardware compatibility. The most basic compression metric is raw size reduction: parameter count, memory footprint in bytes, and compressed storage requirements. Size alone, however, is misleading. On ImageNet, MobileNetV2 achieves approximately 72 percent top-1 accuracy with 3.5M parameters vs. ResNet-50's 76 percent accuracy with 25.6M parameters, about 7.3× fewer parameters at comparable accuracy, or roughly 6.9× more accuracy per parameter (Sandler et al. 2018; He et al. 2016a).

Pruning benchmarks must distinguish between structured and unstructured approaches, because they produce qualitatively different results on real hardware. Structured pruning removes entire neurons or filters, yielding smaller dense operations that conventional kernels can exploit (H. Li et al. 2017). Unstructured pruning eliminates individual weights and can produce very sparse models, but realizing actual speedups requires specialized sparse computation support—meaning benchmark protocols must specify hardware platform and software implementation (Han et al. 2015; Gale et al. 2019).

Quantization benchmarks evaluate precision reduction across data types. In the illustrative MobileNetV2 scenario, INT8 reduces raw weight storage by 4× and the assumed latency values yield the displayed speedup; realized performance depends on the hardware and implementation. The precision-accuracy trade-off is analyzed in section 12.8.2 and the energy implications in section 12.9.1. Mixed-precision approaches push further by applying different precision levels to different layers: critical layers retain FP16 while computation-heavy layers use INT8 or INT4, enabling fine-grained efficiency optimization. Knowledge distillation adds another dimension: a smaller student model can preserve much of a teacher's behavior while reducing size and inference cost, but benchmarking must verify that the student generalizes rather than merely memorizing the teacher's outputs (Hinton et al. 2015).

Critically, acceleration factors vary dramatically across hardware platforms: sparse models, reduced-precision models, and efficient architectures only deliver speedups when the target runtime has kernels, memory layouts, and accelerator support that exploit them. Current benchmark suites like MLPerf focus primarily on standardized reference models, while production deployments often use compressed or hardware-specific variants. This gap between what benchmarks measure and what production actually runs remains one of the field's most consequential blind spots.

12.5.11 Mobile and edge benchmarks

Mobile and edge deployments face constraints radically different from cloud environments, requiring specialized benchmarking approaches that capture the unique trade-offs in resource-constrained settings. These constraints form an interdependent triangle of power consumption, inference latency, and model accuracy, where improving any two typically degrades the third. Edge deployment requires navigating trade-offs that cloud deployments can largely ignore, summarized in table 12.6.

Table 12.6: Edge vs. Cloud Deployment Constraints: Typical cloud deployments treat power as an operating cost and latency as a service target, while edge deployments often face tighter energy and response-time limits that constrain achievable accuracy.

Constraint	Cloud Impact	Edge Impact
Power	Operational cost	Device energy budget
Latency	Service-level target	Local response deadline
Accuracy	Task-quality target	Balanced with power/latency

As a concrete example, a smartphone camera AI for real-time object detection may need to process video-rate inputs while staying inside a tight thermal envelope. In that setting, a MobileNet-family model can be the correct benchmark target even if a larger ResNet-family model reports higher accuracy in a cloud setting, because the edge benchmark must include sustained latency, power, and thermal behavior. A sustained edge benchmark exposes these gaps between marketed specifications

and operational behavior. The peak-versus-sustained gap established in section 12.1 turns acute at the edge for a physical reason absent in the data center: a passively cooled device cannot shed the heat of continuous inference indefinitely, so burst-mode numbers can degrade under thermal throttling. That thermal mechanism, not measurement sloppiness, makes edge benchmarking a categorically different exercise than cloud benchmarking.

Example 12.3: Benchmarking the edge

Scenario: An edge engineering team benchmarks a smart doorbell chip with a vendor-advertised “1 W AI” power specification.


Diagnosis: Early burst-mode testing matches vendor claims, but continuous inference loop execution builds up junction heat, triggering thermal throttling and dropping steady-state throughput.

Systems lesson: Short burst benchmarks obscure long-term thermal throttling. Benchmarking edge ML workloads requires continuous, sustained execution testing to reflect true operational hardware limits.

Systems Perspective 12.7: Edge benchmark reality check

When evaluating edge hardware claims, four factors determine whether vendor numbers translate to real-world performance:

1. **Peak vs. sustained:** A vendor may advertise 45 TOPS peak throughput while a sustained thermal run delivers closer to 20 TOPS. Always benchmark under sustained workloads longer than 30 s.
2. **Power at idle vs. active:** In this scenario, a device consuming 50 mW idle and 2 W active could report active draw for marketing, but if the application runs inference 1 percent of the time, effective power draw is ~69.5 mW, not 2 W.
3. **Thermal envelope:** Edge devices often operate inside a narrow TDP envelope. Exceeding it triggers throttling, so benchmark reports omitting thermal conditions are incomplete.
4. **End-to-end vs. accelerator-only:** NPU benchmarks often exclude data transfer overhead. Moving image data from camera to NPU and back can exceed inference time for small models.



Burst benchmark FPS collapses once thermal throttling sets in.

Thermal throttling in a constrained passive-cooling envelope can begin during sustained inference, making short burst benchmarks misleading for always-on applications. Any edge evaluation must therefore account for sustained power draw under thermal steady state, not burst-mode peaks, and must measure end-to-end latency including data transfer overhead.

12.5.11.1 Heterogeneous processor coordination

Mobile SoCs integrate heterogeneous processors (CPU, GPU, DSP, NPU) requiring specialized benchmarking that captures workload distribution complexity while accounting for thermal and battery constraints. Effective processor coordination can deliver large gains when work is placed on the processor that matches its compute pattern. Each processor excels at different workload profiles: CPUs handle control flow, small batches, and sequential processing; GPUs accelerate parallel floating-point operations and general ML inference; DSPs excel at fixed-point signal processing and always-on detection tasks; and NPUs target specific neural network architectures with INT8/INT4 precision.

Benchmarks must evaluate workload placement decisions, not just individual processor performance. A voice assistant, for example, might use a low-power DSP for always-on wake-word detection, switch to an NPU for a short speech-recognition burst, and use the CPU for language understanding. Single-processor benchmarks miss these orchestration dynamics entirely.

12.5.11.2 Battery and thermal benchmarking

Battery impact varies dramatically by use case: computational photography can consume watts during active capture, while background AI for activity recognition may need to stay in a milliwatt-scale budget for acceptable all-day endurance. The challenge is that *instantaneous power draw* during inference tells only part of the story; what matters for battery life is the *total energy budget* across a realistic usage pattern.

The most important factor is the workload duty cycle: what fraction of time the system actually runs inference. A doorbell camera that processes occasional frames spends nearly all its time idle, making standby power the dominant concern. A real-time video analytics pipeline, by contrast, is inference-bound almost continuously, making per-inference energy the critical metric. Background power, the energy consumed when the model is loaded but waiting for input, bridges these extremes and often exceeds inference energy for intermittent workloads. Finally, sustained thermal behavior must be characterized over minutes rather than seconds, because edge devices that deliver impressive burst performance frequently throttle as junction temperatures rise, settling at substantially lower steady-state throughput.

12.5.11.3 Edge-cloud coordination

Mobile benchmarking must also evaluate 5G/Wi-Fi edge-cloud coordination, with URLLC¹⁹ emphasizing very low latency and high reliability for critical applications. This coordination introduces benchmarking dimensions absent from purely local evaluation. Network latency variability means that inference pipelines splitting work between device and cloud face unpredictable round-trip costs. Fallback behavior determines what happens when connectivity fails entirely: whether the device degrades gracefully to a smaller on-device model or queues requests until connectivity resumes. Workload splitting decisions (what computation runs locally vs. remotely) and privacy constraints (what data can be transmitted for cloud inference) further shape the benchmark design space. Each of these dimensions must be measured under realistic network conditions rather than idealized lab connectivity.

Automotive deployments add Automotive Safety Integrity Level (ASIL) validation, multi-sensor fusion, and wide-temperature environmental testing. These unique requirements necessitate comprehensive frameworks evaluating sustained performance under thermal constraints, battery efficiency across usage patterns, and connectivity-dependent behavior, extending beyond isolated peak measurements.

Whether benchmarking cloud servers or microcontrollers, however, a critical distinction cuts across all deployment contexts: the same neural network behaves entirely differently depending on whether it is *learning* or *predicting*. This distinction shapes what is measured, how it is measured, and which metrics matter—and it is so fundamental that separate benchmarking frameworks have emerged for each phase.

12.6 Training vs. Inference

The same accelerator can fail in opposite ways: a training job may waste days because gradient synchronization dominates, while an inference service may miss its SLO because tail latency spikes under bursty traffic. Training and inference therefore create evaluation requirements so different that separate benchmarking frameworks emerged for each: MLPerf Training and MLPerf Inference (Mattson et al. 2020; Reddi et al. 2019). The critical question is whether theoretical TFLOP/s translate to practical time-to-train or queries-per-second. Training seeks optimal parameters through iterative refinement (chapter 8), processing billions of examples over hours or days, stressing memory bandwidth, multi-GPU scaling, and sustained throughput. Inference applies those parameters to individual inputs in serving systems (chapter 13), often within millisecond deadlines, stressing latency consistency, cold-start time (model startup delay), and power efficiency; chapter 14 connects those measurements to rollout and monitoring practice.

The differences cascade through every aspect of system design. Training involves forward and backward passes, while inference normally performs forward passes with fixed parameters. Memory allocation diverges sharply because training requires parameters, gradients, optimizer states, and saved activations, often creating several-times-higher demand than weight-only inference. The factor depends on the optimizer, numerical precision, activation checkpointing, batch and sequence shape,

¹⁹ | URLLC (ultra-reliable low-latency communication): ITU-R defines a 1 ms radio-interface user-plane latency target and 99.999 percent successful delivery within 1 ms for a 32-byte packet under specified test conditions. These radio-interface targets are not a sub-1-ms end-to-end application guarantee. Edge-inference benchmarks should report radio or network latency and compute latency separately and together, alongside model quality.

and inference caches. Training employs mixed-precision computation and gradient compression to manage this overhead, while inference uses more aggressive precision reduction (detailed in section 12.8.2) and techniques like post-training quantization and knowledge distillation. Resource utilization patterns also contrast: training targets sustained GPU saturation, whereas inference contends with variable request patterns that can leave hardware underutilized, as the roofline analysis in section 12.3.2 demonstrated.

Energy costs follow different patterns. Training energy costs are amortized across model lifetime and measured in total energy per trained model; estimates for large training runs can reach the scale of thousands of megawatt-hours (GPT-3 has been estimated at roughly 1,287 MWh) (Patterson et al. 2021). Inference energy costs accumulate per query and can become a dominant operational consideration at scale. A durable way to reason about per-query energy is the identity $E_{\text{total}} = \text{Power} \times T$. For example, if measured average accelerator power during the inference window is 300 W, a 10 ms inference uses $300\text{W} \times 0.01\text{s} = 3\text{J}$, or about 0.0008 Wh; at 100 ms, that becomes about 0.0083 Wh. A TDP rating is not a substitute for this runtime power measurement.

The training-vs.-inference distinction guides benchmark design by highlighting which metrics matter most for each phase and how evaluation methodologies must differ. Training benchmarks emphasize convergence time and scaling efficiency; inference benchmarks prioritize latency consistency and resource efficiency across diverse deployment scenarios. Training benchmarks come first because the quality of the trained model sets the ceiling for everything inference can deliver.

12.7 Training Benchmarks

In an illustrative procurement failure, a team purchases a larger GPU cluster expecting proportional training-speed gains, only to discover that communication overhead and memory bottlenecks limit the actual speedup. **Training benchmarks** exist to catch this kind of gap before procurement. They divide into three categories: convergence metrics that measure learning progress, throughput metrics that measure computational efficiency, and scalability metrics that measure distributed performance.

Definition 12.3: ML training benchmarks

ML training benchmarks are machine learning system benchmarks that measure the time to reach a target quality metric (for example, a specified validation accuracy or loss threshold) on a fixed dataset and model, quantifying the rate of convergence per unit of resource.

1. **Significance:** Training benchmarks reveal large gaps invisible to hardware specs. Holding the model and quality target fixed, the time to convergence can vary widely across hardware-software stacks because training performance depends on the full pipeline: data loading (D_{vol}/BW), compute utilization (η_{hw}), gradient synchronization (L_{lat}), and fault recovery overhead. A peak FLOP/s spec sheet captures none of these interactions.
2. **Distinction:** Unlike inference benchmarks, which measure per-query latency and throughput under load, training benchmarks measure time-to-accuracy across the full optimization loop: data loading, forward pass, backward pass, gradient synchronization, and optimizer step. The binding constraint shifts from compute (R_{peak}) at small scale to communication (BW) at large scale.
3. **Common pitfall:** A frequent misconception is that training benchmarks measure “how fast the GPU runs.” At large scale, interconnect bandwidth (BW) for gradient synchronization and fault tolerance overhead (checkpoint I/O, straggler mitigation) often dominate the benchmark result more than peak FLOP/s.

Training benchmarks validate whether hardware acceleration delivers promised training throughput. The GPU clusters, TPU pods, and distributed training strategies examined in chapter 11 all claim dramatic speedups, and training benchmarks reveal which claims hold under realistic workloads. They evaluate how hardware configurations, data loading mechanisms, and distributed training strategies perform when training production-scale models. These benchmarks are vital because training represents the largest capital expenditure in ML systems, and only rigorous time-to-accuracy

²⁰ | **GPT-3 (Generative Pre-trained Transformer 3):** OpenAI's 2020 language model (175B parameters, 300B training tokens) consumed an estimated 3,640 petaFLOP-days on 10,000 V100 GPUs (Patterson et al. 2021). This scale illustrates why training benchmarks are essential for predicting whether a planned training run is operationally viable before committing the compute.

measurement reveals whether that capital delivers proportional value rather than dissipating into scaling inefficiencies, memory bottlenecks, or communication overhead.

For instance, large-scale models like OpenAI's GPT-3²⁰ (Brown et al. 2020), which consists of 175B parameters trained on approximately 570 GB of filtered CommonCrawl text (from a ~45 TB raw dataset, combined with other sources to form 300B training tokens), highlight the immense computational demands of modern training. Standardized ML training benchmarks provide systematic evaluation of the underlying systems to ensure that hardware and software configurations can meet these unprecedented demands efficiently.

12.7.1 Training benchmark motivation

MLPerf Training (Mattson et al. 2020; MLCommons 2024c) provides the standardized framework for this kind of time-to-quality measurement. Figure 12.5 shows that performance improvements across successive benchmark versions have outpaced the plotted Moore's Law baseline, with some workloads achieving large multi-year speedups (Tschand et al. 2024). The comparison illustrates a core principle: what gets measured gets improved. The standardized benchmarking framework creates competitive pressure that drives rapid optimization across the entire ML computing stack.




Figure 12.5: MLPerf Training Progress: Workload lines report cumulative relative performance across benchmark versions against the plotted cumulative Moore's Law baseline. Every workload's latest point exceeds the baseline at the same date, and Mask R-CNN reaches the largest plotted gain at 48.8x; the comparison captures combined hardware and software progress, not hardware scaling alone. Source: (Tschand et al. 2024).

Beyond charting that progress, training benchmarks uncover the inefficiencies that systematic evaluation makes visible: slow data loading, underutilized accelerators, excessive memory overhead, and communication bottlenecks that erode scaling efficiency. The theoretical hardware capabilities established in chapter 11 (for example, GPU TFLOP/s, TPU tensor throughput) only translate to actual training speedups when benchmarks verify them under realistic conditions.

Training benchmarks serve four interconnected functions. First, they enable hardware and software optimization by providing vendor-neutral comparisons across accelerator architectures and frameworks (TensorFlow, PyTorch) on standardized tasks, guiding hardware selection for data centers and cloud environments. Software optimizations including mixed-precision training²¹ and memory-efficient data loading are similarly quantified. Second, they evaluate scalability: adding GPUs should reduce training time proportionally, but communication overhead, synchronization latency, and memory bottlenecks limit scaling efficiency in practice. Training benchmarks quantify these losses, revealing whether infrastructure investments deliver proportional returns. Third, they provide cost and energy accountability: with large-scale training runs consuming thousands of megawatt-hours, benchmarks that track cost per training run and power consumption per unit of progress help organizations balance computational power with sustainability goals. Finally, standardized evaluation criteria, controlled randomness, submission rules, and audits make comparisons more reproducible and constrain implementation-specific shortcuts.

²¹ | **Mixed-precision training:** Uses lower precision for most arithmetic while preserving higher-precision accumulation where needed (Micikevicius et al. 2017). The benchmarking consequence: mixed-precision and full-precision runs are not directly comparable because reduced memory traffic and larger feasible batch sizes can change convergence dynamics. MLPerf addresses this by fixing the accuracy target, making time-to-accuracy the comparable quantity regardless of precision strategy.

12.7.2 Training metrics

From a systems perspective, training benchmarks assess how efficiently a model reaches a predefined accuracy threshold. Metrics like throughput and scalability are only meaningful relative to whether the model achieves its target accuracy; without this constraint, optimizing raw speed may be misleading. MLPerf Training codifies this by defining specific accuracy targets per task: a system that trains quickly but misses the target is invalid, and one that converges accurately but too slowly is impractical. Effective benchmarking balances speed, efficiency, and accuracy convergence.

12.7.2.1 Time and throughput

One of the primary metrics for evaluating training efficiency is the time required to reach a predefined accuracy threshold. Training time (T_{train}) measures how long a model takes to converge to an acceptable performance level, reflecting the overall computational efficiency of the system. Let $\text{Accuracy}(t)$ be the model's accuracy at training time t , and let **target accuracy** be the benchmark-specific threshold (for example, 75.9 percent top-1 accuracy for ResNet-50 on ImageNet in MLPerf). Equation 12.1 formally defines this metric, keeping the benchmark focused on how quickly a system achieves meaningful results:

$$T_{\text{train}} = \inf \{t \geq 0 : \text{Accuracy}(t) \geq \text{target accuracy}\} \quad (12.1)$$

A run that never reaches the target has no finite time-to-accuracy, preventing raw training speed from rewarding a system that fails the benchmark's quality criterion.

Throughput,²² often expressed as the number of training samples processed per second, provides an additional measure of system performance. Let N_{samples} be the total number of training samples processed and T_{train} the training time from equation 12.1. Equation 12.2 makes the rate explicit by dividing processed samples by training time:

$$\text{Throughput} = \frac{N_{\text{samples}}}{T_{\text{train}}} \quad (12.2)$$

Throughput alone does not guarantee meaningful results, as a model may process a large number of samples quickly without necessarily reaching the desired accuracy. For example, MLPerf Training specifies workload-specific quality targets; a ResNet-50 result on ImageNet must reach a top-1 accuracy target of 75.9 percent to be valid (Mattson et al. 2020; MLCommons 2024c). A hypothetical system that processes many images per second but fails to reach the target is not a valid benchmark result, while a slower system that converges efficiently can be preferable. This highlights why throughput should be evaluated in relation to time-to-accuracy rather than as an independent performance measure.

12.7.2.2 Scalability and parallelism

Scalability measures how effectively training performance improves as resources are added. Ideally, doubling GPU count should halve training time. In practice, communication overhead, memory bandwidth limits, and parallelization inefficiencies constrain scaling well below linear.

Napkin Math 12.3: Scaling efficiency calculation

Problem: A team trains ResNet-50 on ImageNet. Single-GPU training takes 24 hours. With 8 GPUs, training takes 4 hours. Is this good scaling? Where did the efficiency go?

Step 1: Define scaling efficiency. For strong scaling (fixed problem size, more processors), let $T(1)$ be the training time on a single GPU, $T(N_{\text{GPU}})$ the training time on N_{GPU} GPUs, and N_{GPU} the GPU count. Equation 12.3 defines efficiency:

$$\text{Eff}_{\text{scaling}} = \frac{T(1)}{N_{\text{GPU}} \times T(N_{\text{GPU}})} \times 100\% \quad (12.3)$$

Step 2: Calculate efficiency. $\text{Eff}_{\text{scaling}}(8) = \frac{24 \text{ hours}}{8 \times 4 \text{ hours}} \times 100\% = 24/32 = 75 \text{ percent}$

²² | **Throughput:** Originating in manufacturing to measure units produced per unit time, the term entered computing in the 1960s batch-processing era. The manufacturing origin carries a systems lesson: throughput and latency are inherently opposed, because batching increases throughput (more units per hour) at the cost of individual item wait time. In ML serving, this manifests as the batch-size trade-off: larger batches improve GPU utilization but increase per-request latency.

With perfect scaling, 8 GPUs would complete in 3 hours (24 hours/8 GPUs). The actual 4 hours represents 75 percent efficiency.

Step 3: Account for the efficiency loss. Table 12.7 decomposes the “missing” 25 percent into measurable overhead categories—gradient synchronization, memory copy, load imbalance, and batch-size effects—each measurable through a distinct profiling signal.

Table 12.7: Scaling Efficiency Loss Sources: Illustrative decomposition of the missing efficiency budget into measurable overhead categories and the corresponding measurement signal used to attribute each loss. Actual percentages depend on model, interconnect, storage, batch schedule, and framework implementation.

Source	Example Contribution	Measurement
Gradient synchronization	10–15%	AllReduce time per step
Memory copy (CPU↔GPU)	3–5%	Data transfer profiling
Load imbalance	2–5%	Per-GPU step time variance
Batch size effects	2–5%	Larger batches converge differently

Step 4: The systems insight. Scaling efficiency decreases as N_{GPU} grows because communication overhead scales with GPU count while per-GPU compute shrinks. In this worked example, eight GPUs reach 75 percent efficiency; at larger scales, the same arithmetic makes clear why sophisticated communication and input-pipeline optimization become necessary.

MLPerf reports raw performance, while scaling efficiency can be calculated across matched submissions: a system achieving $2\times$ throughput at 50 percent efficiency may be worse than $1.5\times$ throughput at 90 percent efficiency, depending on cost constraints.

When training large-scale models such as GPT-3, OpenAI employed a large cluster of NVIDIA V100 GPUs in a distributed training setup (Brown et al. 2020; Patterson et al. 2021). Google’s TPU v4 systems demonstrate the same distributed-systems lesson at data center scale: adding computational resources provides more raw power, but performance and resiliency depend on network communication, topology, and operational management (Jouppi et al. 2023; Zu et al. 2024). Benchmarks such as MLPerf quantify how well a system scales across multiple accelerators, providing insights into where inefficiencies arise in distributed training.

Parallelism in training is categorized into data parallelism, model parallelism, and pipeline parallelism (see chapter 8), each presenting distinct challenges. Data parallelism, the most commonly used strategy, involves splitting the training dataset across multiple compute nodes. The efficiency of this approach depends on synchronization mechanisms and gradient communication overhead. In contrast, model parallelism partitions the neural network itself, requiring efficient coordination between processors. Benchmarks evaluate how well a system manages these parallelism strategies without degrading accuracy convergence. A key metric for evaluating parallelism is **scaling efficiency**, which quantifies how much of the added computational capacity translates into actual speedup.

12.7.2.3 Resource utilization

The efficiency of machine learning training depends not only on speed and scalability but also on how well available hardware resources are used. Compute utilization measures the extent to which processing units, such as GPUs or TPUs, are actively engaged during training. Low utilization may indicate bottlenecks in data movement, memory access, or inefficient workload scheduling.

For instance, when training BERT on a TPU cluster, input-pipeline inefficiencies can limit overall throughput even when the accelerators have high raw compute power. If storage retrieval or preprocessing cannot keep up, the system fails to keep the TPUs fully busy. Profiling resource utilization identifies the bottleneck, and optimizations such as prefetching, caching, and more parallel input processing can improve sustained performance.

Memory bandwidth is another critical factor, as deep learning models require frequent access to large volumes of data during training. If memory bandwidth becomes a limiting factor, increasing compute power alone will not improve training speed. Benchmarks assess how well models use

available memory, ensuring that data transfer rates between storage, main memory, and processing units do not become performance bottlenecks.

I/O performance also plays a direct role in training efficiency, particularly when working with large datasets that cannot fit entirely in memory. Benchmarks evaluate the efficiency of data loading pipelines, including preprocessing operations, caching mechanisms, and storage retrieval speeds. Systems that fail to optimize data loading can experience large slowdowns, regardless of computational power.

12.7.2.4 Energy efficiency and cost

Training large-scale machine learning models requires substantial computational resources, leading to considerable energy consumption and financial costs. Energy efficiency metrics quantify the power usage of training workloads, helping identify systems that optimize computational efficiency while minimizing energy waste. The increasing focus on sustainability has led to the inclusion of energy-based benchmarks, such as those in MLPerf Training, which measure power consumption per training run. The same power accounting governs inference, where precision becomes the dominant energy lever; section 12.9 works through why INT8 quantization cuts per-inference energy by attacking both memory traffic and arithmetic cost.

Training GPT-3 was estimated to consume 1,287 MWh of electricity (Patterson et al. 2021). If a system can achieve the same accuracy with fewer training iterations, it directly reduces energy consumption. Energy-aware benchmarks help guide the development of hardware and training strategies that optimize power efficiency while maintaining accuracy targets.

Cost considerations extend beyond electricity usage to include hardware expenses, cloud computing costs, and infrastructure maintenance. Training benchmarks provide insights into the cost-effectiveness of different hardware and software configurations by measuring training time in relation to resource expenditure. Organizations can use these benchmarks to balance performance and budget constraints when selecting training infrastructure.

12.7.2.5 Fault tolerance and robustness

Training workloads often run for extended periods, sometimes spanning days or weeks, making fault tolerance an essential consideration. A resilient system must handle unexpected failures (hardware malfunctions, network disruptions, and memory errors) without compromising accuracy convergence.

In large-scale cloud-based training, node failures are an operational reality. If a GPU node in a distributed cluster fails, training must continue without corrupting the model. Production systems checkpoint for fault tolerance, periodically saving progress so a failure does not restart the run. For large language model (LLM) training, however, checkpointing is itself a systems bottleneck: a single checkpoint must write model weights plus optimizer states to network storage, which at 100-billion-parameter scale can mean hundreds of gigabytes written before training resumes. During that write, accelerators can stall, degrading time-to-accuracy by extending the effective iteration time. Production LLM training systems address this by overlapping checkpoint I/O with the next training step (asynchronous checkpointing) or by using high-bandwidth parallel file systems that reduce idle time. MLPerf Training itself primarily measures time-to-quality under standardized workloads and does not benchmark failure recovery directly, but checkpoint overhead is a material component of any real sustained-throughput number.

12.7.2.6 Reproducibility and standardization

Reproducibility studies have repeatedly shown that modest benchmark gains can disappear when random seeds, hardware, framework versions, or implementation details change (Henderson et al. 2018). This failure mode illustrates a pervasive problem: training benchmarks involve stochastic processes (weight initialization, data shuffling, dropout masks) that interact with hardware-specific behaviors (floating-point rounding, memory layout, compiler optimizations) to produce results that can vary meaningfully across environments.

A deeper layer of non-determinism comes from the parallel hardware itself. Operations such as parallel atomic additions, used during gradient accumulation for sparse embeddings in models like Graph Neural Networks, execute in non-deterministic order across threads when concurrent

updates target the same memory location. The resulting floating-point summation order changes across runs, producing bit-for-bit different gradients even with identical inputs and seeds. Enforcing bit-exact reproducibility in these cases requires disabling the parallel accumulation paths, which reduces training throughput—a direct trade-off between reproducibility and performance that benchmark protocols must explicitly address. Without explicit controls for all these sources of variability, benchmark numbers reflect a specific confluence of conditions rather than a system’s genuine capability.

MLPerf Training addresses this through standardized data preprocessing, target-quality rules, and repeated accepted runs that characterize stochastic variation (Mattson et al. 2020). The point is not merely to produce a fast run, but to show that the reported performance reflects system capability rather than a favorable combination of stochastic factors.

For a training benchmark, reproducibility is therefore the full run envelope, not just the random seed. A credible report must preserve the model commit, dataset checksum, preprocessing pipeline, seed plan, framework and compiler versions, precision policy, batch schedule, hardware topology, thermal and power limits, and checkpoint behavior. It must also report the distribution of accepted runs rather than a single best run. Only then can the benchmark separate a real system improvement from a favorable interaction among software version, hardware state, and stochastic training path.

12.7.3 Training performance evaluation

A comprehensive training benchmark considers multiple dimensions of system behavior because each dimension identifies a different way hardware investment can fail to become convergence. Table 12.8 summarizes the core categories and associated metrics commonly used to benchmark system-level training performance, providing a framework for understanding how training systems behave under different workloads and configurations.

Table 12.8: Training Benchmark Dimensions: Key categories and metrics for evaluating machine learning training systems beyond simple speed, covering resource efficiency, reproducibility, and overall performance trade-offs across different training approaches and infrastructure configurations.

Category	Key Metrics	Example Benchmark Use
Training Time and Throughput	Time-to-accuracy (seconds, minutes, hours); Throughput (samples/sec)	Comparing training speed across different GPU architectures
Scalability and Parallelism	Scaling efficiency (percent of ideal speedup); Communication overhead (latency, bandwidth)	Analyzing distributed training performance for large models
Resource Utilization	Compute utilization (percent GPU/TPU usage); Memory bandwidth (GB/s); I/O efficiency (data loading speed)	Optimizing data pipelines to improve GPU utilization
Energy Efficiency and Cost	Energy consumption per run (MWh, kWh); Training throughput per watt (FLOP/s/W)	Evaluating energy-efficient training strategies
Fault Tolerance and Robustness	Checkpoint overhead (time per save); Recovery success rate (percent)	Assessing failure recovery in cloud-based training systems
Reproducibility and Standardization	Variance across runs (percent difference in accuracy, training time); Framework consistency (TensorFlow vs. PyTorch vs. JAX)	Ensuring consistency in benchmark results across hardware

These dimensions interact in ways that tables cannot capture. Higher throughput from reduced precision (for example, TF32) is meaningless if it increases the iterations required to reach target accuracy, making time-to-accuracy the essential corrective metric. Scaling efficiency can look nearly linear at small node counts but taper as gradient synchronization costs dominate. Resource utilization metrics reveal why: a BERT pretraining task with moderate GPU utilization may be bottlenecked by its data pipeline, not its accelerators. Checkpointing for fault tolerance introduces its own overhead, requiring balance between resilience and performance.

Across all dimensions, measurement accuracy depends on controlling for hardware variability. GPU boost clock²³ behavior and thermal throttling²⁴ can shift results enough to swamp small claimed gains, making repeated runs and statistical rigor (as established earlier) essential for distinguishing genuine performance differences from noise.

23 | GPU boost clock: Dynamic frequency scaling raises clocks above base when thermal and power headroom permit. The benchmarking trap: short benchmark runs can capture boost-clock performance, but sustained ML training may settle to lower steady-state frequencies as junction temperature rises. Reporting burst-phase results overstates the throughput a production workload can sustain.

24 | Thermal throttling: Frequency reduction triggered when junction temperature exceeds safe limits. For edge devices without active cooling, throttling can begin during sustained inference, meaning peak throughput numbers from short benchmarks may misrepresent steady-state performance.

Despite the availability of well-defined benchmarking methodologies, misleading conclusions recur when teams treat one training metric as a substitute for the whole optimization loop. The following pitfalls show where the benchmark must keep speed, convergence, scaling, and reproducibility tied together.

Training benchmark failures usually start when throughput is treated as the objective rather than as one part of the learning process. A system can increase examples per second by using lower numerical precision, reducing synchronization, or even bypassing certain computations, but those changes only help if convergence is preserved. A TF32 run may outpace FP32 per step and still lose overall if numerical instability increases the number of iterations required to reach the target accuracy. The benchmark therefore has to report throughput in relation to time-to-accuracy, ensuring that speed optimizations do not come at the expense of convergence efficiency.

Scaling creates a second trap because a small-node result can look linear until communication and synchronization dominate. The earlier eight-GPU calculation shows why small-node results cannot be extrapolated linearly once synchronization becomes the binding term.

As the scaling efficiency calculation in section 12.7.2.2 demonstrated (where 8 GPUs achieved only 75 percent efficiency), extrapolating single-node results to clusters is a common error. Google's experience with 4,096-node TPU v4 clusters shows this effect at extreme scale, where synchronization challenges become the dominant performance factor. Proper benchmarking should measure scaling efficiency explicitly rather than assuming linear improvement.

The same discipline applies to failures and interference. Many benchmarks assume idealized conditions where hardware failures, network instability, and workload interference do not occur, even though those events are routine at scale. Effective benchmarking accounts for checkpointing overhead, failure recovery efficiency, and resource contention rather than reporting only best-case performance.

Reproducibility poses another threat. Results must reproduce across stacks: a TensorFlow run with Accelerated Linear Algebra optimizations may exhibit different convergence behavior than the same model trained in PyTorch with Automatic Mixed Precision (AMP), because floating-point arithmetic, memory layouts, and optimization strategies can all shift training time and accuracy.

Avoiding these pitfalls requires evaluating throughput in relation to accuracy convergence, assessing scaling efficiency holistically, and accounting for real-world failures rather than assuming idealized conditions. A model trained efficiently, however, still requires validation of its deployment performance, which shifts the evaluation framework entirely.

12.8 Inference Benchmarks

Training benchmarks measure how quickly a system learns; inference benchmarks measure how reliably it serves. This shift changes nearly every aspect of evaluation. Training tolerates variable iteration times as long as convergence proceeds; inference requires consistent latency because users experience every slow response. Training optimizes for aggregate throughput across hours; inference must handle unpredictable request patterns and scenario-specific deadlines. Large benchmark training commonly uses dedicated high-performance hardware, whereas inference spans environments from data center GPUs to mobile phones to microcontrollers.

This is where the optimization chapters converge: the accelerated hardware from chapter 11 runs compressed models from chapter 10 to deliver real-time predictions. Inference benchmarks reveal whether those theoretical speedups become actual latency reductions under realistic deployment conditions.



Definition 12.4: ML inference benchmarks

ML inference benchmarks are machine learning system benchmarks that quantify a system's ability to meet latency constraints (L_{lat}) at specified throughput levels, measuring scenario-defined latency statistics, throughput (queries per second), and power efficiency across representative serving scenarios.

1. **Significance:** Inference benchmarks expose the workload- and system-specific gap between unconstrained throughput and throughput while meeting a service-level objective

(SLO), such as a p99 latency target. Queuing delays can push tail latency above the target at high load, a gap that is invisible without a benchmark that enforces latency targets at each throughput level.

2. **Distinction:** Unlike training benchmarks, which measure time-to-accuracy over a fixed dataset, inference benchmarks measure per-query response time under realistic load patterns, capturing queuing effects, batching trade-offs, and cold-start overhead that determine real-world serving economics.
3. **Common pitfall:** A frequent misconception is that average latency is a sufficient benchmark. A system with low average latency but a long p99 tail can violate a percentile-based production SLO for the slowest 1 percent of requests; at high request rates, that small percentage becomes a large number of affected users. The benchmark must report the latency statistic named by the service objective rather than assuming the mean or p99 is universally decisive.

12.8.1 Inference benchmark motivation

Large-scale benchmark training commonly runs on dedicated data center hardware, whereas inference spans dramatically different deployment scenarios—from real-time applications like autonomous driving and conversational AI to mobile devices, IoT systems, and embedded processors. This diversity extends to hardware: while GPUs and TPUs dominate large-scale training, inference workloads often use specialized accelerators like NPUs, field-programmable gate arrays, and dedicated inference chips such as Google’s Edge TPU.²⁵ Inference benchmarks evaluate how well hardware selection, model optimization, and data pipeline design work together across these deployment environments.

Scaling inference workloads across cloud servers, edge platforms, mobile devices, and TinyML systems introduces additional complexity. Figure 12.6 reveals the staggering power consumption differentials among these systems—spanning over ten orders of magnitude from microwatts in tiny embedded devices to hundreds of kilowatts in data center training clusters. The ranges are representative rather than exhaustive. This spread explains why no single benchmark can serve all deployment contexts: a metric meaningful for data center optimization (kilowatts per rack) becomes irrelevant for battery-powered edge devices (milliwatts per inference). Inference benchmarks must evaluate the trade-offs between latency, cost, and energy efficiency within each scale to assist organizations in making informed deployment decisions.

²⁵ | **Edge TPU (tensor processing unit):** Google’s fixed-function edge AI accelerator. It illustrates a benchmarking constraint specific to fixed-function accelerators: its headline throughput applies only to quantized TensorFlow Lite models with supported operator types, so models requiring unsupported operators fall back to the host CPU or need graph rewrites before the accelerator result is meaningful.




Figure 12.6: Power Consumption Differentials: Power usage spans over ten orders of magnitude across ML system types, from microwatts in TinyML devices through watts at the edge to kilowatts in data center inference and hundreds of kilowatts for training clusters. Ranges are representative and vary by hardware and workload.

These deployment differences create the practical motivation for inference benchmarks: they evaluate the bottlenecks that emerge when models transition from development to production serving. The motivating factors parallel those for training (hardware optimization, scalability, cost, fair comparison) but differ in specifics. Software optimization frameworks apply inference-specific techniques such as operator fusion (see chapter 10 and chapter 11), precision calibration, and kernel tuning, whose impact on latency, throughput, and power efficiency must be measured under realistic conditions to confirm they deliver real improvements without degrading accuracy. Auto-tuning compilers add a hidden variable: the compiler itself can require hours of optimization per model-hardware pair, meaning benchmark results reflect the tuning budget as much as the hardware capability, and comparing results across submissions requires normalizing for compiler optimization time.

Scalability concerns also shift character. Training scales by adding GPUs to reduce time-to-accuracy on a fixed workload, whereas inference must scale dynamically in response to fluctuating user demand, handling traffic spikes without violating latency guarantees. Cold-start performance, the time required for a model to load and begin processing queries, becomes a distinct inference concern with no training analog. Applications that load models on demand, such as serverless AI deployments, are particularly sensitive to this overhead.

The cost and energy profile of inference differs sharply from training. Training concentrates cost in discrete runs that may recur as data, objectives, or models change; inference cost accumulates continuously with production traffic. Running an inefficient model at scale can multiply cloud compute expenses, and on battery-powered devices, excessive computation directly affects usability. Benchmarks that measure cost per inference request and efficiency per watt help organizations optimize for both performance and sustainability across deployment platforms.

MLPerf Inference extends the standardized comparison principles established for training benchmarks to deployment scenarios, defining evaluation criteria for tasks such as image classification, object detection, and speech recognition across different hardware platforms. This ensures that inference performance comparisons remain meaningful and reproducible while accounting for deployment-specific constraints like latency requirements and energy efficiency (Reddi et al. 2019).

12.8.2 Inference metrics

For example, a voice assistant must respond quickly enough that users do not perceive lag, while a recommendation engine must score enough candidates to keep pace with user scrolling. These constraints (latency and throughput) define the performance envelope within which all serving optimizations must operate. Inference metrics formalize these real-world demands into measurable quantities, and they differ from training metrics in kind, not just degree, because the optimization target shifts from “how fast can we learn?” to “how reliably can we serve?” Training cares about throughput and time-to-accuracy; inference cares about latency consistency, resource efficiency, and deployment practicality, spanning cloud data centers handling millions of requests to edge devices operating under strict power constraints.

12.8.2.1 Latency and tail latency

Latency (introduced in chapter 2) measures the time for an inference system to process an input and produce a prediction. Average latency is useful, but it does not capture high-percentile delays that degrade reliability in high-demand scenarios.

To account for this, benchmarks often measure tail latency,²⁶ which summarizes the high-latency tail rather than the worst case. These values are commonly reported as the 95th percentile (p95) or 99th percentile (p99) latency, meaning that 95 percent or 99 percent of inferences are completed within a given time. Applications with hard deadlines require explicit deadline-miss or worst-case analysis beyond these percentiles.

These measurements form the basis for Service Level Objectives (SLOs) and SLAs, which formalize performance expectations.

²⁶ | **Tail latency:** A high-percentile response time, such as p95 or p99, determines production SLO or SLA compliance only when that objective or agreement specifies the percentile. Dean and Barroso (2013) showed that in fan-out architectures (common in recommendation systems), even 1 percent slow responses compound: a request touching 100 backend shards has about a 63 percent chance that at least one shard hits its 1 percent tail. Benchmarks reporting only mean latency hide this failure mode.



Definition 12.5: SLOs and SLAs

SLOs and SLAs are performance commitment specifications for production ML serving systems: a service-level objective (SLO) is a target value or range for a service-level indicator, while an SLA is a contract that specifies service objectives and the consequences of failing to meet them.

1. **Significance:** A latency SLO constrains the L_{lat} term in the iron law by setting a target for a specified latency indicator. A representative production setup might set the internal SLO tighter than the external SLA, leaving operational headroom for transient spikes, maintenance windows, and cascading failures.
2. **Distinction:** An SLO has no contractual consequence by itself, while missing an SLA objective triggers the consequences specified in the agreement, which may include credits or penalties. Teams commonly set an internal SLO tighter than the external agreement to leave operational headroom, but this is a practice rather than a definitional requirement.
3. **Common pitfall:** A frequent misconception is that meeting average latency satisfies a tail-latency SLO. SLOs can use averages, percentiles, rates, or other indicators; when the objective is defined at p99 or p99.9, an excellent mean does not establish compliance.

The distinction matters in practice: engineering teams optimize toward SLOs while the business commits to SLAs. Choosing the wrong metric to optimize wastes engineering effort or violates customer guarantees.

A usable latency objective also fixes the request population and measurement window. Otherwise, two teams can report the same percentile threshold over different traffic mixes and reach incompatible conclusions about compliance.



Checkpoint 12.2: Metric selection

The metric shapes the optimization.

Apply three rules before finalizing metric selection:

- ☐ Throughput vs. latency: Decide whether the workload optimizes for cost or user experience, then name the metric the benchmark must hold fixed.
- ☐ Tail latency: State what p99 exposes that mean latency hides.
- ☐ End-to-end: List the preprocessing, model execution, and postprocessing stages included in the reported latency.

Tail latency's connection to user experience at scale becomes critical in production systems serving millions of users. Even small p99 latency degradations create compounding effects across large request volumes: if 1 percent of requests experience $10\times$ latency (for example, 1000 ms instead of 100 ms), this affects 10,000 requests per million, potentially leading to timeout errors, poor user experience, and customer churn. Search engines and recommendation systems demonstrate this sensitivity: Google's search-latency experiments found measurable reductions in daily searches per user after 100–400 ms server-side delays (Brutlag 2009), which is why interactive services often treat sub-100 ms response times as a practical design target.

Service level objectives (SLOs) in production systems therefore focus on tail latency rather than mean latency to ensure consistent user experience. Interactive services often define percentile-based latency objectives because occasional slow responses have disproportionate impact on user satisfaction. Large-scale systems may track even deeper tails, such as p99.9, when traffic spikes and infrastructure variation affect reliability.

The challenge of meeting these tail latency targets is that the source of the tail is often architectural, not algorithmic. A garbage-collected runtime, a shared kernel driver, or a priority-inversion bug in the serving stack can inject latency spikes that no model optimization will remove.

📖 War Story 12.1: Discord Read States rewrite (2019)

Context: Discord, a real-time chat platform supporting millions of concurrent users, originally implemented its Read States service in Go, tracking which channels and messages each user read (Howarth 2020).

Mechanism: Go's forced garbage collection passes scanned an LRU cache holding tens of millions of entries, causing periodic "stop-the-world" GC pauses every two minutes regardless of memory tuning.

Impact: Tail latency (P_{99}) spiked dramatically every two minutes, degrading user experience across millions of active connections.

Fix: In 2019, Discord rewrote the Read States service in Rust, eliminating garbage collection pauses and dropping average response times to microseconds.

Systems lesson: Mean latency alone does not describe user experience when the service objective is defined at the tail. Language-runtime choices can set a floor on that tail: feature stores that serve real-time embeddings for recommendation models may use managed runtimes, and garbage-collection pauses can delay every downstream inference request waiting for retrieval. Model optimization cannot close a gap whose bottleneck lies in the retrieval path. The Discord incident documents the underlying mechanism outside ML and shows why the complete serving stack belongs inside a production latency benchmark.

12.8.2.2 End-to-end vs. component latency

A critical distinction in inference benchmarking is between component latency (time spent in model computation) and end-to-end latency (total time from request arrival to response delivery). Many benchmarks report only model inference time, obscuring the remaining overhead that determines actual user experience. The overhead is not marginal: serialization, network hops, and queue wait time can dominate total request time, making model-only optimizations yield diminishing returns.

📖 Example 12.4: The JSON serialization trap

Scenario: Researchers evaluating the Clipper low-latency model serving system decomposed end-to-end request latency across network, CPU serialization, and model execution (Crankshaw et al. 2017).

Diagnosis: For lightweight models, CPU-side JSON deserialization and IPC data copying consume more wall-clock time than accelerator neural network execution.

Systems lesson: Reporting isolated model inference latency misses end-to-end serving bottlenecks. Production benchmarking must account for CPU data parsing, IPC serialization, and network transport overhead.

Table 12.9 gives an illustrative latency breakdown for an inference request. The model inference stage that vendors report as their "benchmark" number spans 5 to 100 ms, yet the queue wait time it sits behind ranges from 0 to over 1,000 ms: under load, the single component a benchmark measures is dwarfed by one it never sees, so the reported number can be a small slice of what the user actually experiences.

Table 12.9: Inference Latency Breakdown: Different pipeline components contribute to end-to-end latency, and model inference (the number vendors typically report) can be only one part of total request time. Queue wait time can dominate under load, making end-to-end measurement essential for realistic performance assessment. Ranges are illustrative rather than universal.

Component	Example Range	Notes
Network round-trip	10–100 ms	Varies by region
Request parsing	0.1–1 ms	JSON/protobuf
Input preprocessing	1–50 ms	Tokenization, image resize
Queue wait time	0–1000+ ms	Load-dependent

Component	Example Range	Notes
Model inference	5–100 ms	The “benchmark”
Output postprocessing	0.5–10 ms	Decoding, format
Response serialization	0.1–1 ms	JSON/protobuf

These component-level contributions explain why optimizing any single stage yields diminishing returns on end-to-end performance, an *optimization ceiling* formalized by Amdahl’s Law.

Napkin Math 12.4: Amdahl’s Law: Optimization ceiling

Problem: A vision pipeline spends 8 ms on preprocessing (JPEG decode, resize, normalize) and 10 ms on inference. If inference alone is optimized by 5×, how much does end-to-end latency actually improve?

Math: Optimizing inference from 10 ms to 2 ms reduces total latency from 18 ms to only 10 ms, a 1.8× improvement rather than 5×. Amdahl’s Law formalizes this ceiling: if preprocessing consumes fraction f of total latency, then even infinitely fast inference yields at most $1/f$ speedup. With preprocessing at 44.4 percent of total latency, identifying the dominant fraction $f = 0.444$ yields a maximum achievable speedup $1/f$ of 2.25×, regardless of model optimization.

Systems insight: Aggressive model optimization yields disappointing end-to-end results whenever the nonmodel fraction dominates. Any component speedup quoted in isolation shrinks once the untouched stages are measured alongside it, and the shortfall grows as those stages take a larger share of the pipeline. Comprehensive benchmarks must either include preprocessing in measurements or state explicitly that reported speedups apply only to the inference component.

Amdahl’s ceiling highlights why rigorous benchmarking methodology matters. Comprehensive latency reporting requires specifying which components are included, measuring under realistic load conditions, and distinguishing component from end-to-end metrics. Before interpreting any benchmark result, verify that the measurement approach itself is sound.

Throughput and batch efficiency measure whether a serving system can use available hardware without violating latency constraints. Throughput counts how many inference requests a system processes per second, typically expressed as queries per second (QPS) or frames per second (FPS). Single-instance systems process each input independently on arrival; batch systems process multiple inputs in parallel, exploiting hardware parallelism for higher efficiency.

Checkpoint 12.3: Benchmarking methodology

Bad benchmarks optimize the wrong things.

Three practices distinguish rigorous benchmarks from misleading ones:

- ☐ Representative data: Compare the benchmark input distribution against the production distribution.
- ☐ Warm-up: State how many initial runs were discarded before recording measurements.
- ☐ Isolation: Document the machine, competing workloads, and runtime configuration used during measurement.

For example, cloud-based services handling millions of queries per second benefit from batch inference, where large groups of inputs are processed together to maximize computational efficiency. In contrast, applications like robotics, interactive AI, and augmented reality require low-latency single-instance inference, where the system must respond immediately to each new input. Benchmarks must consider both single-instance and **batch throughput** to provide a comprehensive understanding of inference performance across different deployment scenarios.

Speed alone is insufficient because inference optimizations can change model behavior. Reducing numerical precision can accelerate computation while cutting memory and energy, as the illustrative MobileNetV2 energy estimate shows (table 12.17), but lower-precision calculations can introduce accuracy degradation. Inference benchmarks therefore evaluate how well models perform under different numerical settings, such as FP32, FP16, and INT8.²⁷ Many modern AI accelerators support mixed-precision inference, allowing systems to use numerical representations selected for workload requirements. Model compression techniques²⁸ further improve efficiency, but their impact on model accuracy varies depending on the task and dataset. Benchmarks help determine whether these optimizations are viable for deployment, ensuring that improvements in efficiency do not come at the cost of unacceptable accuracy loss.

Memory footprint and model load time define whether the model can start, stay resident, and respond within the deployment envelope. Unlike training, where models can span multiple accelerators, inference often runs within strict memory budgets. Total model size determines storage requirements, RAM usage reflects working memory during execution, and memory bandwidth can bottleneck data transfer between processing units. Cold-start performance becomes critical when models are loaded on demand rather than kept resident in memory. In serverless AI environments,²⁹ where resources scale dynamically with incoming requests, the time from idle to active execution determines whether users experience acceptable response times.

Model load time refers to the duration required to load a trained model into memory before it can process inputs. In some cases, particularly on resource-limited devices, models must be reloaded frequently to free up memory for other applications. The time taken for the first inference request is also an important consideration, as it reflects the total delay users experience when interacting with an AI-powered service. Benchmarks help quantify these delays, ensuring that inference systems can meet real-world responsiveness requirements.

Deployment-scale metrics extend the same logic from one request to a workload. Cloud services must handle millions of concurrent users efficiently, allocating resources dynamically as demand fluctuates without compromising latency; mobile devices must manage multiple simultaneous AI models without overloading the system. Scalability measures how well inference performance improves when additional computational resources are allocated. In some cases, adding more GPUs or TPUs increases throughput proportionally, but in other scenarios, bottlenecks such as memory bandwidth limitations or network latency may limit scaling efficiency. Benchmarks also assess how well a system balances multiple concurrent models in real-world deployment, where different AI-powered features may need to run at the same time without interference.

Energy consumption closes the loop because inference workloads run continuously in production. Mobile and edge devices face the most acute constraints, where battery life and thermal limits restrict available computational resources. Even in large-scale cloud environments, power efficiency directly impacts operational costs and sustainability goals. The energy required for a single inference is often measured in joules per inference, reflecting how efficiently a system processes inputs while minimizing power draw. In cloud-based inference, efficiency is commonly expressed as queries per second per watt (QPS/W) to quantify how well a system balances performance and energy consumption. For mobile AI applications, optimizing inference power consumption extends battery life and allows models to run efficiently on resource-constrained devices. Reducing energy use also plays a key role in making large-scale AI systems more environmentally sustainable, ensuring that computational advancements align with energy-conscious deployment strategies.

12.8.3 Inference performance evaluation

Unlike training, inference systems must process inputs and deliver predictions efficiently across diverse deployment scenarios. Latency, throughput, memory usage, and energy efficiency provide the structured measures for evaluating this performance.

Table 12.10 should be read as a deployment filter: each metric identifies a constraint that can dominate a different serving environment. Percentile latency can characterize interactive serving tails, while safety-critical systems with hard deadlines require a deadline-miss probability or worst-case bound appropriate to the safety requirement. Queries per second per watt governs a battery-bound mobile deployment, where the same model is judged on endurance rather than peak speed. Trade-

27 | INT8 (8-bit integer):

INT8 sits at the aggressive end of the precision hierarchy (FP32 baseline, FP16 halves raw weight storage, INT8 quarters it), and each step demands increasing care to preserve accuracy. Post-training INT8 quantization normally uses a representative calibration dataset; accuracy depends on the model, operator coverage, quantization method, and how well calibration data represent deployment inputs. INT8 benchmarks must specify whether they use post-training quantization or quantization-aware training and document any calibration procedure.

28 | Model compression

benchmarking: Compression impact must be measured across four dimensions simultaneously: accuracy degradation, inference speedup, memory reduction, and energy savings. A technique achieving 10× size reduction with 1 percent accuracy loss may still be unsuitable if latency does not improve proportionally; unstructured pruning, for example, reduces parameter count but rarely improves latency on dense hardware because sparse operations lack efficient hardware support on most GPUs.

29 | Serverless AI:

Deployment paradigm where models scale from zero instances on demand. Cold-start time varies with the model, runtime, storage path, hardware allocation, and provider. Benchmarks for intermittent workloads must report whether initialization and model loading are included, because warm-instance latency alone can understate user-perceived latency.

offs between metrics, including speed vs. accuracy and throughput vs. power consumption, are common, and understanding these trade-offs is essential for effective system design.

Table 12.10: Inference Performance Metrics: Latency, throughput, and resource usage metrics provide a quantitative basis for optimizing deployed machine learning systems and selecting appropriate hardware configurations, balancing speed, cost, and accuracy in production applications.

Category	Key Metrics	Example Benchmark Use
Latency and Tail Latency	Mean latency (ms/request); Tail latency (p95, p99, p99.9)	Evaluating interactive and soft real-time performance
Throughput and Efficiency	Queries per second (QPS); Frames per second (FPS); Batch throughput	Comparing large-scale cloud inference systems
Numerical Precision Impact	Accuracy degradation (FP32 vs. INT8); Speedup from reduced precision	Balancing accuracy vs. efficiency in optimized inference
Memory Footprint	Model size (MB/GB); RAM usage (MB); Memory bandwidth utilization	Assessing feasibility for edge and mobile deployments
Cold-Start and Load Time	Model load time (s); First inference latency (s)	Evaluating responsiveness in serverless AI
Scalability	Efficiency under load; Multi-model serving performance	Measuring robustness for dynamic, high-demand systems
Power and Energy Efficiency	Power consumption (W); Performance per W (QPS/W)	Optimizing energy use for mobile and sustainable AI

These metrics interact through unavoidable trade-offs. Optimizing for high throughput via large batch sizes increases latency, making a system unsuitable for real-time applications. Reducing numerical precision improves power efficiency and speed but may degrade accuracy. The deployment environment determines which trade-offs are acceptable: cloud systems prioritize scalability and throughput, while edge devices are dominated by memory and power constraints. Evaluating inference performance holistically, rather than fixating on a single metric, ensures that systems meet their functional, resource, and performance goals in context.

Deployment scenario determines the priority order among those metrics. The operational constraints and success criteria vary dramatically across contexts, so metric priorities help engineers focus benchmarking effort and interpret results within the right decision framework. Table 12.11 illustrates how performance priorities shift across five major deployment contexts, revealing the systematic relationship between operational constraints and optimization targets.

Table 12.11: Illustrative Performance Metric Priorities by Deployment Context: A sample priority ordering shows how operational constraints can change benchmark selection and result interpretation; actual thresholds must come from the deployment specification.

Deployment Context	Primary Priority	Secondary Priority	Tertiary Priority	Key Design Constraint
Real-Time Applications	Latency	Reliability	Memory Footprint	User experience demands immediate response
Cloud-Scale Services	Throughput (QPS)	Cost Efficiency	Average Latency	Business viability requires massive scale
Edge/Mobile Devices	Power Consumption	Memory Footprint	Latency	Battery life and resource limits dominate
Training Workloads	Training Time	GPU Utilization	Memory Efficiency	Research velocity enables faster experimentation
Scientific/Medical	Accuracy	Reliability	Explainability	Correctness cannot be compromised for performance

The key insight is that the same metric can be primary in one context and irrelevant in another. Latency ranks first for real-time applications (autonomous vehicles must process sensor data within strict timing deadlines) but tertiary for cloud services (which accept higher latency in exchange for cost efficiency per query). A smartphone AI assistant that improves throughput while increasing power consumption may regress when battery life is the binding constraint. A medical diagnostic system may rationally prefer a slower model with higher validated accuracy. This context-dependence

means that a $2\times$ throughput improvement represents substantial value for cloud deployments but minimal benefit for battery-powered edge devices, where 20 percent power reduction delivers superior operational impact.

Even with well-defined metrics, inference evaluations fail when the benchmark ignores the deployment constraint that dominates the serving system. The following pitfalls show where average latency, memory, energy, cold starts, and scaling assumptions can each invalidate an otherwise plausible result.

Inference benchmark failures begin when the benchmark averages away the event users actually notice. Tail latency (p95, p99) determines production reliability, not mean latency; a conversational AI system that misses its tail-latency target will produce unacceptable response delays even if its average response time looks healthy. Resource constraints create the same kind of mismatch. A model with excellent cloud throughput may still be unusable on a phone or edge device if its memory footprint or power draw exceeds the deployment budget, so practical inference benchmarks must include memory and energy alongside latency.

Serverless and on-demand serving add a separate first-request constraint. **Cold-start latency**³⁰ measures the time required to initialize a model and process the first request, so excluding model load time creates unrealistic expectations for responsiveness. Evaluating both model load time and first-inference latency ensures that systems are designed for the conditions they will actually face.

Inference benchmarks also become misleading when one metric is optimized in isolation. Maximizing batch throughput can degrade latency, while aggressive precision reduction can reduce accuracy. A precision example makes the comparability problem concrete.

Numerical precision optimization exemplifies this challenge particularly well. Accelerators can provide substantially higher INT8 operation throughput³¹ than FP32 floating-point throughput, creating compelling performance narratives. Those narratives are only valid when the benchmark also checks accuracy, supported operator coverage, and whether the reported operations are comparable across devices.

Scaling and application fit require the same skepticism. The linear scaling pitfall discussed for training benchmarks applies equally to inference, though the bottlenecks differ: training scaling is often limited by gradient synchronization, while inference scaling encounters memory bandwidth saturation, thermal throttling under sustained load, and request-routing overhead, the extra time spent assigning requests to model replicas in distributed serving. These limitations arise from physical hardware constraints and interconnect architectures (chapter 11). A cloud-optimized benchmark can therefore be irrelevant for an edge deployment where energy and memory dominate, so benchmark selection has to follow the application requirement rather than the most convenient leaderboard.

Finally, inference results need the same statistical discipline as training results. MLPerf uses sufficiently long runs and large query counts, and its latency metric depends on the scenario: SingleStream reports a 90th-percentile latency estimate, Server enforces a 99th-percentile latency constraint while maximizing throughput, and Offline reports throughput (Reddi et al. 2019). These protocols capture serving behavior that a mean or a single timing measurement would miss.

12.8.4 MLPerf inference benchmarks

Avoiding these pitfalls requires treating inference benchmarking as a process of balancing multiple priorities (latency, throughput, memory, energy, and accuracy) rather than optimizing for any single metric in isolation; MLPerf Inference operationalizes that balance through deployment-specific scenarios. MLPerf Inference matters because deployment context changes what a result means. The benchmark, developed by MLCommons,³² provides a standardized framework for evaluating machine learning inference performance across a range of deployment environments. MLPerf began with training benchmarks in 2018; MLPerf Inference was added later to standardize deployment-time evaluation across scenarios. As machine learning systems expanded into diverse applications, it became clear that a one-size-fits-all inference benchmark was insufficient. The resulting family of MLPerf inference benchmarks maps each benchmark to a deployment setting, so a score can be interpreted against the latency, throughput, memory, and power constraints the system will face.

³⁰ | **Cold-start latency:** The initialization time from idle state includes storage reads, host-to-accelerator transfer, and framework setup. For a 7B-parameter model in FP16 (~14 GB) already resident in host memory, PCIe 4.0 transfer at 25 GB/s effective bandwidth alone takes ~560 ms; storage and initialization add to it. This lower bound makes cold-start mitigation (model caching, speculative loading) a systems design requirement.

³¹ | **TOPS (tera operations per second):** A measure of raw computational throughput (trillions of operations/second). The H100 delivers 1979 TOPS INT8 vs. the Apple M2 Neural Engine at 15.8 TOPS and Edge TPU at 4 TOPS, but these numbers conflate different operation types—multiply-accumulate vs. accumulate vs. activation. TOPS comparisons across vendors are meaningful only when the operation definition, precision, and sparsity assumptions are identical, conditions rarely met in vendor specifications.

³² | **MLCommons:** Launched in 2020 as a nonprofit consortium evolving from the 2018 MLPerf effort, MLCommons includes members across industry and academia (MLCommons 2026a). Requiring published system specifications improves result comparability, but it does not prevent submitters from selecting which systems or workloads to enter. Published results reveal large performance differences between vendors on identical workloads, making MLCommons the closest the field has to SPEC-style apples-to-apples hardware comparison.

12.8.4.1 MLPerf Inference

MLPerf Inference (Reddi et al. 2019) serves as the baseline inference benchmark, defining standardized scenarios for deployment-time evaluation across data-center and edge settings. Submissions are split into two tracks: the **closed division** enforces strict model structure and precision equivalence to ensure direct hardware-to-hardware comparison (apples-to-apples), whereas the **open division** allows submitters to alter model architecture, quantization algorithms, or retrain models to demonstrate creative algorithmic co-design. It assesses performance across deep learning workloads such as image classification, object detection, natural language processing, and recommendation systems. This version of MLPerf is a widely used reference point for comparing AI accelerators, GPUs, TPUs, and CPUs when the submission rules and workload scenario match the intended deployment environment.

Major technology companies regularly reference MLPerf results for hardware procurement decisions. When evaluating hardware for recommendation systems infrastructure, MLPerf benchmark scores on DLRM³³ workloads can inform choices between different accelerator generations. Across generations, benchmark results often show substantial throughput improvements, although the magnitude depends on workload, software stack, and system configuration. This illustrates how standardized benchmarks can translate into consequential infrastructure decisions.

These standardized evaluations provide invaluable comparisons, but the cost of comprehensive benchmarking limits who can participate and how thoroughly systems are evaluated.

Systems Perspective 12.8: The cost of comprehensive benchmarking

Submitting to MLPerf can require dedicated engineering effort, hardware access, tuning, validation, and careful documentation across the relevant configurations. The cost depends heavily on workload and scope, but it is high enough that submissions are dominated by major technology companies and hardware vendors, while smaller organizations rely on published results rather than conducting their own comprehensive evaluations. That barrier motivates the need for more lightweight, internal benchmarking practices that organizations can use to make informed decisions without the expense of full-scale standardized benchmarking.

The rest of the MLPerf inference family narrows that baseline by deployment context. MLPerf Mobile (MLCommons 2024a) evaluates whether a model can remain responsive within smartphone power and memory limits (Janapa Reddi et al. 2022), measuring image classification, object detection, image segmentation, and language-processing workloads. MLPerf Client (MLCommons 2026b) addresses the local-computing decision: whether consumer devices can run AI workloads directly rather than relying on cloud inference. Its current emphasis on local generative-AI and LLM workloads makes CPUs, discrete GPUs, and integrated NPUs part of the benchmarked system rather than incidental host hardware. MLPerf Tiny (C. Banbury et al. 2021) tests the extreme constraint case: embedded and ultra-low-power AI systems, such as IoT devices, wearables, and microcontrollers. These variants preserve the same benchmark discipline while changing the binding resource from data center throughput to client responsiveness, mobile power, or microcontroller memory.

12.8.4.2 MLPerf execution scenarios

The same hardware can report dramatically different benchmark numbers depending on *how* requests arrive—a fact that explains why vendor claims often fail to predict production performance. Classic MLPerf Inference defines four execution scenarios that characterize distinct traffic patterns, each requiring different optimization strategies (Reddi et al. 2019). Current client and generative-AI benchmark variants also include interactive measurements for latency-sensitive LLM workloads, where metrics such as time-to-first-token and time-per-output-token become central (MLCommons 2026b).

SingleStream SingleStream processes one request at a time, measuring latency for sequential inference. This scenario models mobile and embedded applications where a single user interacts with the device: a smartphone camera app classifying images, a voice assistant processing speech, or a wearable detecting gestures. The key metric is per-request latency, and batching provides no

³³ | **DLRM (deep learning recommendation model):** Facebook’s 2019 recommendation architecture combines embedding tables for categorical features with multilayer perceptrons for continuous features (Naumov et al. 2019). DLRM stresses benchmarks differently than vision or language models: its embedding tables can be large enough that memory capacity and bandwidth dominate compute throughput. That makes DLRM a useful memory-bound recommendation workload in MLPerf-style inference evaluation, revealing hardware limitations invisible to compute-bound benchmarks (Reddi et al. 2019).

benefit since requests arrive only after the previous result is consumed. Optimization focuses on preprocessing efficiency and power consumption rather than throughput.

MultiStream MultiStream processes multiple synchronized input streams simultaneously, modeling scenarios like autonomous vehicles with multiple cameras that must be processed together for spatial fusion. Unlike SingleStream’s sequential requests, MultiStream requires processing frames from all sensors within tight video-rate deadlines. The key distinction from Server mode is that MultiStream inputs arrive in lockstep, while Server requests arrive independently and unpredictably. The key constraint is synchronization: all streams must complete before the planning module can act. Optimization focuses on jitter handling and meeting hard deadlines rather than average throughput.

Server Server generates requests following a Poisson distribution, simulating cloud API traffic where requests arrive independently and unpredictably. This scenario models web services handling millions of queries from different users. Unlike SingleStream’s guaranteed sequential arrival, Server traffic creates queuing dynamics where multiple requests compete for resources. The key metrics are throughput (queries per second) and tail latency (p99), and dynamic batching can improve efficiency by grouping requests that arrive within a time window. Optimization balances throughput against latency SLOs.

Offline Offline provides all inputs upfront, measuring maximum throughput when latency constraints are removed. This scenario models batch processing pipelines: overnight data processing, scientific computing, or precomputing recommendations. With no latency requirement, systems can use maximum batch sizes to saturate hardware utilization. The key metric is pure throughput (samples per second), and optimization focuses entirely on hardware efficiency.

Table 12.12 maps the classic execution scenarios, plus the newer Interactive LLM-oriented case, to their deployment contexts and optimization strategies.

Table 12.12: MLPerf Execution Scenarios: Classic MLPerf Inference scenarios map to distinct deployment contexts, each requiring different optimization strategies, and newer Interactive scenarios extend the framework to latency-sensitive LLM workloads. SingleStream and MultiStream prioritize latency, Server balances throughput and latency, Offline maximizes throughput, and Interactive emphasizes token-level responsiveness. Matching the scenario to deployment context determines which benchmark results are relevant.

Scenario	Context	Strategy	Focus
SingleStream	Mobile apps, embedded devices	No batching (batch = 1)	Preprocessing, power efficiency
MultiStream	Autonomous driving, video analytics	Synchronized sensor fusion	Jitter handling, deadline guarantees
Server	Cloud APIs, web services	Dynamic batching with timeout	Throughput-latency trade-off tuning
Offline	Batch processing, data pipelines	Maximum batch size	Throughput, hardware utilization
Interactive	Chat, agents, local generative AI	Token streaming, KV-cache management	Time-to-first-token, time-per-output-token

Each scenario acts as a workload contract by fixing which form of latency or throughput a valid result must preserve.



Lighthouse 12.2: MobileNetV2 on EdgeTPU

Extending lighthouse 12.1, the hardware-acceleration claim from chapter 11 is evaluated using an illustrative SingleStream-style edge-inference scenario (Reddi et al. 2019; C. Banbury et al. 2021).

Hardware acceleration claim: In this illustrative edge-accelerator scenario, assume INT8 MobileNetV2 inference takes ~2 ms on the accelerator, approximately 7.5× faster than a Cortex-M-class CPU (~15 ms). Actual results depend on operator coverage, clock frequency, thermal state, and implementation.

Table 12.13 reports the illustrative SingleStream-style scenario assumptions.

Table 12.13: EdgeTPU vs. Cortex-M7 MobileNetV2 Scenario: Illustrative SingleStream-style inference latency, end-to-end latency, power, and energy values show how preprocessing overhead narrows the headline accelerator speedup; these are not controlled measurements.

Metric	CPU (Cortex-M7)	EdgeTPU	Headline Ratio	System-Level Ratio
Inference latency	~15 ms	~2 ms	7.5× faster	—
End-to-end latency	~18 ms	~6 ms	—	~3× faster
Power consumption	~120 mW	~500 mW	—	~4.2× higher
Energy per inference	~1.8 mJ	~1 mJ	—	~1.8× more efficient

What this reveals: Under these assumptions, the 7.5× inference speedup narrows to ~3× end to end because preprocessing runs on the CPU in both cases. EdgeTPU uses more active power but completes faster, yielding lower inference energy; deployment requires controlled measurement.

The deployment decision depends on the workload. For battery-powered devices running infrequently, the active inference calculation alone favors EdgeTPU, but total battery impact depends on sleep power, wake-up energy, host-transfer overhead, and whether the accelerator adds idle leakage while the system waits. For continuous video operation, EdgeTPU’s lower active energy per inference is much more likely to dominate.

The SingleStream result illustrates why benchmarking requires matching the MLPerf scenario to the deployment context: SingleStream emphasizes latency, while Offline benchmarks would give different conclusions optimized for throughput rather than latency.

The scenarios explain why the same hardware can report dramatically different benchmark numbers. In an illustrative comparison, an accelerator with high Offline throughput can sustain much lower Server-mode throughput once p99 latency constraints and queuing overhead are enforced, because Server mode cannot always use maximum batch sizes. When evaluating hardware for a specific application, selecting the appropriate scenario ensures benchmark results predict production performance. The MobileNetV2 result therefore makes scenario selection a deployment requirement rather than a reporting choice.


Training benchmarks measure learning speed; inference benchmarks measure serving speed. Yet both measures share a critical blind spot: they say nothing about *how much energy* the system consumes to achieve that speed. A system that sets throughput records while consuming kilowatts of power may be economically unsustainable or physically impossible to deploy at the edge. Completing the evaluation picture requires power measurement: measuring the energy cost of performance.

12.9 Power Measurement Techniques

A chip vendor advertises “10 TOPS at 0.5 W,” but under sustained inference load, thermal throttling drops actual throughput to 3 TOPS at 2 W. Without standardized power measurement, this 13.3× efficiency gap between the datasheet and reality goes undetected until deployment.

This third dimension is critical because chapter 11 established TOPS/W as a primary design objective alongside raw TOPS. Power benchmarks validate whether efficiency-optimized accelerators deliver their promised energy savings. TOPS/W is particularly susceptible to gaming precisely because it is a ratio of two separately quotable peaks: a vendor can read the numerator (operations) at the batch size and precision that maximize throughput and the denominator (watts) at a near-idle operating point, so the advertised efficiency describes a state the chip never occupies under real load. Power benchmarks close that loophole by fixing the workload and the measurement window, forcing the numerator and denominator to be read at the same operating point.

However, measuring power consumption in machine learning systems presents challenges distinct from measuring time or throughput. Power varies with temperature, workload phase, and system configuration in ways that performance metrics do not. Table 12.14 quantifies how energy demands of ML models vary dramatically across deployment environments, spanning multiple orders of magnitude from TinyML devices consuming mere microwatts to data center racks requiring kilo-



watts. This wide spectrum illustrates the central challenge in creating standardized benchmarking methodologies (Henderson et al. 2020).

Table 12.14: Power Consumption Spectrum: The representative deployment points listed here span nearly eight orders of magnitude in power demands, from microwatt-scale TinyML devices (150 μ W) to kilowatt-scale ML server racks (10 kW); figure 12.6 extends the same picture down to 5.6 μ W at the TinyML floor and up to roughly 498 kW at the training-cluster ceiling, covering nearly eleven orders of magnitude. This enormous range explains why no single measurement technique or efficiency metric applies universally: benchmarking a 150 μ W neural processor requires fundamentally different instrumentation than measuring a 10 kW server rack.

Category	Device Type	Power Consumption
Tiny	Neural Decision Processor (NDP)	150 μ W
Tiny	M7 Microcontroller	25 mW
Mobile	Raspberry Pi 4	3.5 W
Mobile	Smartphone	4 W
Edge	Smart Camera	10–15 W
Edge	Edge Server	65–95 W
Cloud	ML Server Node	300–500 W
Cloud	ML Server Rack	4–10 kW

This range spans nearly eight orders and requires scale-specific measurement: microwatt-level TinyML demands different instrumentation than kilowatt-scale racks. A comprehensive framework must maintain consistency, fairness, and reproducibility across both.

12.9.1 Power measurement boundaries

Addressing these measurement challenges requires understanding how power consumption is measured at different system scales, from TinyML devices to full-scale data center inference nodes. Figure 12.7 lays out the distinct measurement boundaries for each scenario: components in green fall inside the energy accounting boundary, while components with red dashed outlines are explicitly excluded from power measurements. This distinction matters because where the boundary is drawn determines what counts as “efficient.”




Figure 12.7: Power Measurement Boundaries: MLPerf defines system boundaries for power measurement, ranging from single-chip devices to full data center nodes, to enable fair comparisons of energy efficiency across diverse hardware platforms. These boundaries delineate which components’ power consumption is included in reported metrics, impacting the interpretation of performance results. Source: (Tschand et al. 2024).

Figure 12.7 is organized into three categories, Tiny, Inference, and Training examples, each reflecting different measurement scopes based on system architecture and deployment environment. In TinyML systems, the entire low-power SoC, including compute, memory, and basic interconnects, typically falls within the measurement boundary. Inference nodes introduce more complexity, incorporating multiple SoCs, local storage, accelerators, and memory, while often excluding remote storage and off-chip components. Training deployments span multiple racks, where only selected elements, including compute nodes and network switches, are measured, while storage systems, cooling infrastructure, and parts of the interconnect fabric are often excluded.

The measurement boundary determines which energy is counted; workload composition determines where that energy is spent. TensorFlow Mobile measurements attribute 57.3 percent of inference energy to data movement (Boroumand et al. 2018), showing why component-only measurements and raw operation counts can miss the dominant system cost. The next example therefore decomposes MobileNetV2 energy into arithmetic and memory access.

Napkin Math 12.5: Why INT8 saves energy

Problem: Under the stated 45 nm component-energy model, how much does replacing FP32 with INT8 reduce MobileNetV2 weight-read and arithmetic energy?

Recall from chapter 11 that moving data costs far more energy than computing on it (the energy-movement invariant formalized in chapter 4 and quantified by Horowitz’s energy estimates (Horowitz 2014)). Understanding *why* quantization reduces energy consumption requires decomposing energy into its physical sources. Two dominant factors determine inference energy: compute operations and memory access.

Narrower datatypes generally require less switching and storage energy per operation, so table 12.15 reveals an 18× gap between FP32 and INT8 multiply cost:

Table 12.15: Per-Operation Energy by Precision: Energy cost of a single multiply at FP32, FP16, and INT8, with relative cost normalized to FP32, illustrating why narrower datatypes reduce arithmetic energy. Values follow the 45 nm estimates in Horowitz (2014) and section D.1.

Precision	Multiplier Energy	Relative Cost
FP32	~3.7 pJ/FLOP	1×
FP16	~1.1 pJ/FLOP	0.3×
INT8	~0.2 pJ/FLOP	0.05×

An 8-bit multiplier uses ~18× less energy than a 32-bit floating-point multiplier in this 45 nm component model because narrower arithmetic reduces switching and storage work (Horowitz 2014). Section D.1 catalogs the per-operation estimates behind these ratios; absolute values and ratios vary with circuit design and process technology.

Table 12.16 extends the picture to memory access, with energy cost per byte across each tier of the hierarchy:

Table 12.16: Memory Access Energy by Tier: Per-byte energy cost across the register-to-DRAM hierarchy, with relative cost normalized to a register read, exposing why off-chip memory traffic dominates inference energy budgets. Values follow the same canonical energy-estimate model used in Horowitz (2014) and section D.1.

Memory Level	Energy per Byte	Relative Cost
Register	~0.1 pJ/byte	1×
L1 Cache	~1 pJ/byte	10×
L2 Cache	~5 pJ/byte	50×
DRAM	~160 pJ/byte	1,600×

Memory access dominates: reading one byte from DRAM costs over 1,600× more energy than a register access.

Math: Component energy follows $E_{\text{load}} = \text{model bytes} \times \text{DRAM energy per byte}$ and $E_{\text{compute}} = \text{FLOPs} \times \text{operation energy}$. The FP32 terms sum as $2243 \mu\text{J} + 2,220 \mu\text{J} = 4,463 \mu\text{J}$; the INT8 terms sum as $561 \mu\text{J} + 120 \mu\text{J} = 681 \mu\text{J}$. Dividing the totals gives a 6.6× reduction under this simplified model.

Table 12.17 combines the two effects in a deliberately simplified MobileNetV2 model. It counts one DRAM read per weight and charges every cataloged FLOP at the listed multiplier energy, while excluding activation traffic, cache behavior, additions as a distinct cost, control overhead, and static power:

Table 12.17: MobileNetV2 INT8 Simplified Energy Estimate: One weight read plus a multiplier-energy charge for every cataloged FLOP under a 45 nm component-energy model. This is not a whole-device inference measurement.

Component	FP32 (14 MB)	INT8 (3.5 MB)	Savings
One weight read from DRAM	2243 μ J	561 μ J	4×
Compute (600 MFLOP)	2,220 μ J	120 μ J	18.5×
Total	4,463 μ J	681 μ J	6.6×

Systems insight: Within this simplified model, the weight read and the arithmetic contribute comparable FP32 energy (2243 μ J and 2,220 μ J), and INT8 reduces both, taking the total from 4,463 μ J to 681 μ J. This explains the physical mechanism behind potential savings, but battery-life or per-inference claims require controlled whole-device power measurements.

Shared infrastructure presents additional challenges. In data centers, resources such as cooling systems and power delivery are shared across workloads, complicating attribution of energy use to specific ML tasks. Cooling alone can account for 20–30 percent of total facility power consumption, making it a major factor in energy efficiency assessments (Barroso et al. 2019). Even at the edge, components like memory and I/O interfaces may serve both ML and non-ML functions, further blurring measurement boundaries.

Within a Transformer forward pass, compute intensity and power can vary across kernels. Feed-forward layers commonly use dense matrix multiplications, while attention can be compute- or memory-bound depending on prefill versus decode, sequence and batch shape, and kernel implementation. Dynamic voltage and frequency scaling (DVFS) can further change power with workload demand (Kim et al. 2008). Power benchmarks therefore need sampling and integration windows that capture kernel-level variation rather than inferring whole-model energy from a single instantaneous reading.

Support infrastructure, especially cooling, is a major component of energy consumption in large-scale deployments. Data centers must maintain operational temperatures, typically between 20–25 °C, to ensure system reliability. Cooling overhead is captured in the power usage effectiveness metric, which ranges from 1.1 in highly efficient facilities to over 2.0 in less optimized ones (Barroso et al. 2019). The interaction between compute workloads and cooling infrastructure creates complex dependencies; for example, power management techniques like DVFS not only reduce direct processor power consumption but also decrease heat generation, creating cascading effects on cooling requirements. Even edge devices require basic thermal management.

12.9.2 Computational efficiency vs. power consumption

The relationship between computational performance and energy efficiency is a central trade-off in modern ML system design. Historically, computations per kilowatt-hour doubled about every 1.5 years, documenting rapid long-run gains in computing efficiency (Koomey et al. 2011). Processor frequency scaling, however, exposes a local trade-off: higher frequency often requires higher voltage, so dynamic power can rise faster than delivered throughput, reflecting the voltage-frequency-power relationship that underlies DVFS and its diminishing returns (Le Sueur and Heiser 2010).

In deployment scenarios with strict energy constraints, particularly battery-powered edge devices and mobile applications, optimizing this performance-energy trade-off becomes essential for practical viability. Model optimization techniques offer promising approaches to achieve better efficiency without material accuracy degradation. Numerical precision optimization techniques, which reduce computational requirements while maintaining model quality, demonstrate this trade-off effectively. Integer quantization studies show that reduced-precision computation can often preserve model quality while improving inference speed, memory traffic, and energy efficiency, although the realized gain depends on model, calibration method, and hardware support (Jacob et al. 2018; Wu et al. 2020; Gholami et al. 2021).

Optimization strategies span three interconnected dimensions: accuracy, computational performance, and energy efficiency. Advanced optimization methods enable fine-tuned control over this trade-off space. Similarly, model optimization and compression techniques require careful

balancing of accuracy losses against efficiency gains. The optimal operating point among these factors depends heavily on deployment requirements and constraints; mobile applications typically prioritize energy efficiency to extend battery life, while cloud-based services might optimize for accuracy even at higher power consumption costs, benefiting from economies of scale and dedicated cooling infrastructure.

Energy efficiency metrics now occupy a central position in AI system evaluation. Power measurement standards such as MLPerf Power (Tschand et al. 2024) provide standardized frameworks for comparing energy efficiency across hardware platforms and deployment scenarios. These standards enable engineers to systematically balance performance, power consumption, and environmental impact when selecting hardware and optimization strategies.

12.9.3 Standardized power measurement

Power measurement techniques like SPEC Power have long served general computing (Lange 2009), but ML workloads expose a fundamental difficulty: instantaneous power consumption during a single inference can shift rapidly between compute-intensive matrix multiplication and memory-stall phases. MLPerf Power formalizes this problem for ML systems by specifying measurement boundaries, instrumentation, and reporting rules across a wide power range (Tschand et al. 2024). This volatility means that any single-point measurement is misleading, and the act of measurement itself (instrumentation overhead, sampling-induced delays) can perturb the very power profile being characterized.

The core challenge is therefore temporal: characterizing a quantity that fluctuates faster than many measurement instruments can sample. Dense matrix operations in transformer layers create short, intense power spikes that require high-frequency sampling to capture accurately, while CNN inference tends toward more consistent power draw amenable to lower sampling rates. The measurement window must also account for ML-specific warm-up periods, where initial inferences consume more power due to cache population and pipeline initialization. Sliding-window averages over repeated inferences smooth these fluctuations into actionable efficiency numbers, but the window size itself becomes a design parameter that can hide or reveal different aspects of the power profile.

Memory access patterns compound the measurement problem because ML systems often spend more energy moving data than computing on it. Recommendation models like DLRM, for example, can consume more energy on memory access than computation—a pattern that traditional compute-focused power measurement misses entirely. Capturing both compute and memory subsystem power consumption requires instrumenting the full data path, not just the processor.

Heterogeneous accelerator configurations introduce further complexity. GPUs, TPUs, and NPUs each maintain independent power management schemes, and modern SoCs dynamically switch between compute resources based on workload characteristics. Accurate system-level measurement requires synchronized power capture across all active compute units—a challenge that scales with system size. Multi-GPU configurations must account for gradient synchronization energy alongside computation, and multi-node deployments add nontrivial network infrastructure power. At the other extreme, edge deployments must capture the energy cost of model updates and data preprocessing alongside inference itself.

Batch size creates a nonlinear relationship with power consumption that single-point measurements cannot characterize. Larger batches improve compute efficiency (better amortization of memory loads) but increase memory pressure and peak power requirements, meaning the most efficient batch size for throughput may differ from the most efficient batch size for energy. Measurement across multiple batch sizes is essential for a complete efficiency profile. System idle states deserve equal attention, particularly for intermittent edge workloads: a wake-word detection TinyML system that actively processes audio for only a small fraction of operating time may be dominated by idle power consumption rather than inference energy. Finally, sustained ML workloads can cause temperature increases that trigger thermal throttling and alter power consumption patterns—an effect particularly acute in edge devices, where thermal constraints limit sustained performance and make extended benchmarking runs essential for realistic characterization.

12.9.4 MLPerf power case study

MLPerf Power (Tschand et al. 2024) turns power measurement from a device-specific reading into a comparable efficiency claim: how many useful inferences a system delivers per watt under a defined boundary. The methodology applies standardized evaluation principles across data center, edge, and tiny inference settings, where the relevant decision changes from rack operating cost to battery life to microwatt-scale endurance.

Boundary-aware standardization matters because the same hardware family can look efficient or wasteful depending on boundary and workload. By adapting the protocol to CPUs, accelerators, and heterogeneous systems while preserving measurement integrity, MLPerf Power makes cross-platform comparisons meaningful across different computing scales.

The benchmark has accumulated many reproducible measurements submitted by industry organizations, demonstrating submitted hardware capabilities and the sector-wide focus on energy-efficient AI technology. The data-center panel in figure 12.8 shows how normalized energy efficiency has evolved across successive MLPerf Inference versions. The gains differ by workload and start version: RetinaNet and ResNet show the largest plotted increases, while GPT-J, DLRM-v2, and Llama 2 cover fewer benchmark rounds, making their gains less directly comparable.




Figure 12.8: Data-Center Energy Efficiency Gains: Each series is normalized to its own first plotted version, so compare improvement within a workload rather than absolute efficiency across workloads. Source: (Tschand et al. 2024).

RetinaNet and ResNet show the largest plotted gains, each reaching $3\times$ its initial normalized efficiency. BERT and GPT-J reach about $1.9\times$, while DLRM-v2 reaches about $1.7\times$, Llama 2 about $1.5\times$, and RNN-T about $1.3\times$.

Timing protocols and power instrumentation provide the raw data for benchmarking. Raw data alone, however, does not guarantee sound conclusions. Converting measurements into meaningful comparisons requires understanding the systematic sources of error, bias, and misalignment that can make even carefully collected benchmark numbers misleading.

12.10 Benchmarking Best Practices

An inference stack that passes a steady-state lab run can still miss latency targets when production traffic arrives in bursts, or when the input mix shifts toward expensive examples. Training throughput, inference latency, and power efficiency each have established measurement protocols validated through MLPerf, but knowing what to measure is insufficient without understanding what benchmarks cannot capture and why this gap has derailed countless deployments.

Benchmarks make simplifying assumptions that enable standardized comparison but can diverge from production reality. Training suites may freeze datasets and tightly specify randomness; production data and retraining conditions evolve. Inference suites often emphasize a defined steady-state scenario; production traffic can be bursty and time-varying. Power benchmarks control thermal conditions; deployed hardware may encounter different ambient temperatures and sustained-load behavior. Four categories of limitations (statistical, deployment-related, system design, and organizational) determine whether benchmark results translate to deployment success.

12.10.1 Statistical and methodological issues

Benchmark results are only as reliable as the measurements that produce them. Three pervasive issues undermine this reliability if left unaddressed.

Incomplete problem coverage represents one of the most pervasive limitations. Many benchmarks, while useful for controlled comparisons, fail to capture the full diversity of real-world applications. Common image classification datasets such as CIFAR-10 (Krizhevsky 2009) contain a limited variety of images. Models that perform well on these datasets may struggle when applied to more complex, real-world scenarios with greater variability in lighting, perspective, and object composition. This gap between benchmark tasks and real-world complexity means strong benchmark performance provides limited guarantees about practical deployment success.

Statistical insignificance arises when benchmark evaluations are conducted on too few data samples or trials, and it is most acute in settings where the evaluation medium itself introduces variance. Large language model evaluation exemplifies this problem: whether scoring a new LLM against a reference using human preference ratings or an LLM-as-judge protocol, the evaluation signal carries high variance because judges respond differently to prompt phrasing, ordering effects, and response length. A reported two-point preference win can disappear entirely across a different judge configuration or prompt template. Rigorous LLM benchmarking therefore requires statistical methods—bootstrap confidence intervals or paired significance tests—applied across enough prompts and response pairings to separate a genuine capability improvement from evaluation noise. Without sufficient trials and diverse input distributions, benchmarking results will mislead: reported differences reflect evaluation noise rather than genuine capability. The statistical confidence intervals around benchmark scores often go unreported, obscuring whether measured differences represent genuine improvements or measurement noise.

Reproducibility represents a major ongoing challenge. Benchmark results can vary measurably depending on factors such as hardware configurations, software versions, and system dependencies. Small differences in compilers, numerical precision, or library updates can lead to inconsistent performance measurements across different environments. To mitigate this issue, MLPerf addresses reproducibility by providing reference implementations, standardized workloads and run rules, and strict submission guidelines. Even with these efforts, achieving true consistency across diverse hardware platforms remains an ongoing challenge. The proliferation of optimization libraries, framework versions, and compiler flags creates a vast configuration space where slight variations produce different results.

12.10.2 Laboratory-to-deployment performance gaps

Statistical rigor ensures that benchmark measurements are accurate. Accurate measurements of the wrong thing, however, still lead to deployment failures. Benchmarks must also align with practical deployment objectives.

Misalignment with real-world goals occurs when benchmarks emphasize metrics such as speed, accuracy, and throughput, while practical AI deployments require balancing multiple objectives including power efficiency, cost, and robustness. A model that achieves top-line accuracy on a benchmark may be impractical for deployment if it consumes excessive energy or requires expensive hardware. Similarly, optimizing for average-case performance on benchmark datasets may neglect tail-latency requirements that determine user experience in production systems. The multi-objective nature of real deployment, encompassing resource constraints, operational costs, maintenance complexity, and business requirements, extends far beyond the single-metric optimization that most benchmarks reward.

12.10.3 System design challenges

Statistical methodology and deployment alignment address how performance is measured and what is optimized. A third category of limitations emerges from the physical systems being measured. Hardware behavior depends on environmental conditions, architectural compatibility, and operational context in ways that complicate fair comparison.

Environmental conditions affect benchmarks in measurable ways. Benchmark results depend on physical conditions (ambient temperature, humidity, altitude) and operational context (background

processes, network load, power supply stability) in subtle but measurable ways. Elevated temperatures trigger thermal throttling that reduces computational speed; background processes compete for resources and alter performance characteristics. Ensuring valid benchmarks requires controlling these factors to the extent possible (temperature-controlled environments, standardized system states, documented background loads) and, when full control is impractical (as in distributed or cloud-based benchmarking), detailed reporting of conditions so that others can account for potential variations when interpreting results.

The hardware lottery³⁴ (Hooker 2021) presents another critical issue. The success of a machine learning model is often dictated not only by its architecture and training data but also by how well it aligns with the underlying hardware. Some models perform exceptionally well not because they are inherently superior but because they map naturally onto GPU or TPU parallel processing capabilities. Other promising architectures may be systematically overlooked because they do not fit dominant hardware platforms.

Hardware compatibility dependence introduces subtle but significant biases into benchmarking results. A model that is highly efficient on a specific GPU may perform poorly on a CPU or a custom AI accelerator. Figure 12.9 makes this hardware dependence concrete by comparing model performance across different platforms. On the CPU uint8 and GPU configurations, the multi-hardware models track the “MobileNetV3 Large min” baseline closely, reaching roughly 77 percent top-1 ImageNet accuracy where the baseline reaches about 75 percent. On the EdgeTPU and DSP hardware the same multi-hardware models sustain that 77 percent at substantially lower latency, while a model tuned only for the CPU would forfeit those gains. This reveals that the “best” model depends entirely on deployment target: a conclusion impossible to reach from single-platform benchmarks.




Figure 12.9: Hardware-Dependent Latency–Accuracy Trade-offs: Each panel holds top-1 ImageNet accuracy on the vertical axis while using a target-specific latency scale horizontally. The arrow links MobileNetV3 Large at roughly equal accuracy on CPU float and DSP, exposing the hardware-dependent latency shift. Source: (Chu et al. 2021).

Without careful benchmarking across diverse hardware configurations, the field risks favoring architectures that “win” the hardware lottery rather than selecting models based on their intrinsic strengths. This bias can shape research directions, influence funding allocation, and impact the design of next-generation AI systems. In extreme cases, it may even stifle innovation by discouraging exploration of alternative architectures that do not align with current hardware trends.

34 | Hardware lottery: Coined by Hooker (2021) to describe how algorithmic success depends on alignment with available hardware and software. Transformer workloads benefit from dense matrix operations that map efficiently to modern accelerators, while graph neural networks and sparse mixture-of-experts models can be harder to evaluate when available silicon and software stacks favor dense kernels. Hardware-specific leaderboards therefore favor architectures aligned with the measured platform and may obscure alternatives that would fare differently under other hardware assumptions.

12.10.4 Organizational and strategic issues

The limitations discussed in this section arise from technical challenges: statistical noise, deployment misalignment, environmental variance, and hardware compatibility. A fourth category emerges from human factors—and these may be the hardest to mitigate because they involve incentives rather than instrumentation. Competitive pressures and research incentives create systematic biases in how benchmarks are used and interpreted. These organizational dynamics require governance mechanisms and community standards to maintain benchmark integrity.

12.10.4.1 Benchmark engineering

While the hardware lottery is an unintended consequence of hardware trends, benchmark engineering is an intentional practice where models or systems are explicitly optimized to excel on specific benchmark tests. This practice can lead to misleading performance claims and results that do not generalize beyond the benchmarking environment.

Benchmark engineering occurs when AI developers fine-tune hyperparameters, preprocessing techniques, or model architectures specifically to maximize benchmark scores rather than improve real-world performance. The distinction between legitimate optimization and benchmark engineering is often blurry, sitting at the threshold where tuning for a specific benchmark crosses into overfitting to it. For example, an object detection model might be carefully optimized to achieve record-low latency on a benchmark but fail when deployed in dynamic, real-world environments with varying lighting, motion blur, and occlusions. Similarly, a language model might be tuned to excel on benchmark datasets but struggle when processing conversational speech with informal phrasing and code-switching.

The pressure to achieve high benchmark scores is often driven by competition, marketing, and research recognition. Benchmarks are frequently used to rank AI models and systems, creating an incentive to optimize specifically for them. While this can drive technical advancements, it also risks prioritizing benchmark-specific optimizations at the expense of broader generalization—precisely the Goodhart’s Law dynamic introduced in section 12.1 and illustrated with the BLEU-score example in section 12.3.1.

12.10.4.2 Bias and over-optimization

The practitioner consuming a benchmark result must determine whether a number reflects legitimate optimization or benchmark engineering. Several practices make that distinguishable, and each catches a specific failure at a specific cost. Transparency is the first line of defense: a submission that documents every optimization applied lets a reader separate general improvement from benchmark-specific tuning, at the cost of exposing techniques a vendor may prefer to keep proprietary. Reporting both benchmark and real-world deployment results closes the same gap from the other side. Diversified evaluation across multiple, continuously updated benchmarks raises the cost of overfitting to any single test set, because a model engineered to win one cannot easily win them all; its cost is the engineering effort of maintaining many benchmarks.

Standardization and third-party verification raise the bar further. Independent audits catch results that fail to reproduce across settings, and the existence proof for this mechanism appears two sections later in MLPerf’s reference-vs-submission validation (section 12.10.5), which disqualifies any submission that cannot hit the reference accuracy target. Application-specific testing catches the failure controlled benchmarks structurally cannot: an autonomous-driving model must be exercised across the weather, lighting, and urban settings it will actually meet, not judged solely on a curated dataset. Multi-hardware testing catches the last case, performance that is really hardware-lottery alignment rather than model quality, by confirming that a result does not depend on compatibility with one platform.

12.10.4.3 Benchmark evolution

A persistent challenge in benchmarking is that benchmarks are rarely static. As AI systems evolve, so must the benchmarks that evaluate them. A performance target that discriminates well under one generation of models, hardware, and applications may lose relevance under another. While benchmarks are essential for tracking progress, they can also become outdated, leading to over-optimization for old metrics rather than real-world performance improvements.

This evolution is evident in the history of AI benchmarks. Early model benchmarks, for instance, focused heavily on image classification and object detection, as these were some of the first widely studied deep learning tasks. However, as AI expanded into natural language processing, recommendation systems, and generative AI, it became clear that these early benchmarks no longer reflected the most important challenges in the field. In response, new benchmarks emerged to measure language understanding (A. Wang et al. 2018, 2019) and generative AI (Liang et al. 2022).

Benchmark evolution extends beyond the addition of new tasks to encompass new dimensions of performance measurement. While traditional AI benchmarks emphasized accuracy and throughput, deployed applications demand evaluation across multiple criteria: fairness, robustness, scalability, and energy efficiency. Figure 12.10 makes these disparate requirements concrete by mapping scientific applications across data rate and computation time. Specifically, Large Hadron Collider sensors must process data at rates approaching 10^{14} bytes per second with nanosecond-scale computation times, while mobile applications operate at 10^4 bytes per second with longer computational windows—a span of roughly ten orders of magnitude in data rate and six to seven in computation time. This range of requirements necessitates specialized benchmarks. For example, edge AI applications benefit from benchmarks like MLPerf that evaluate performance under resource constraints, and scientific application domains need their own “Fast ML for Science” benchmarks (Duarte et al. 2022).




Figure 12.10: Performance Spectrum: Scientific and edge applications occupy widely separated operating points in data rate and available computation time, motivating benchmarks tied to each application’s timing and data-movement constraints. Source: (Duarte et al. 2022).

The need for evolving benchmarks also presents a challenge: stability vs. adaptability. On the one hand, benchmarks must remain stable for long enough to allow meaningful comparisons over time. If benchmarks change too frequently, it becomes difficult to track long-term progress and compare new results with historical performance. On the other hand, failing to update benchmarks leads to stagnation, where models are optimized for outdated tasks rather than advancing the field. Striking the right balance between benchmark longevity and adaptation is an ongoing challenge for the AI community.

Evolving benchmarks remains essential for meaningful progress measurement. Without updates, benchmarks become detached from real-world needs, and researchers optimize for artificial test cases rather than practical challenges. The transition from ImageNet-era accuracy benchmarks to

multi-dimensional evaluations spanning fairness, robustness, and energy efficiency illustrates this evolution in practice.

12.10.5 MLPerf synthesis and benchmark gaming

Benchmark gaming begins when a compiler, runtime, or hardware stack optimizes for the benchmark artifact rather than the workload it is supposed to represent. MLPerf counters that risk by synthesizing the principles discussed throughout this chapter into a single evolving framework: reference implementations and strict submission rules enforce reproducibility, deployment-specific suites (Inference, Mobile, Client, Tiny) align with the three-dimensional evaluation framework, and regular task updates (including generative AI and energy-efficient computing) prevent benchmark stagnation. In the Hennessy & Patterson tradition of quantitative systems, benchmarks are targets, not merely passive measurements. The Goodhart dynamic introduced in section 12.1 applies here in full force. In the high-stakes world of AI hardware, it manifests as benchmark gaming: optimizing hardware or compilers specifically for the benchmark’s unique characteristics, rather than for real-world performance.

Three hypothetical shortcuts illustrate the kinds of benchmark-specific behavior that rigorous run rules must forbid:

- Precision dropping: A compiler silently lowers precision only for the benchmark, without disclosing the changed numerical contract.
- Benchmark-specific operator removal: An implementation exploits the evaluation set or quality threshold to skip work that a general deployment must perform.
- Benchmark detection: A system recognizes the benchmark workload and selects a path unavailable to ordinary inputs.

MLPerf constrains this gaming through reference implementations, quality targets, submission rules, and audits. In the Closed Division, preprocessing, postprocessing, and the model must remain equivalent to the reference, while submitters may change frameworks, layouts, kernels, and numerical representation within those rules. Open submissions may substitute or retrain models but are labeled separately. Accuracy is one guardrail; other rules prohibit benchmark detection and input-dependent optimization.

Yet even the most rigorous system benchmarks validate only one dimension of deployment readiness. A system achieving record throughput and efficiency on MLPerf says nothing about whether the model it runs is accurate on real-world inputs, or whether the data it was trained on represents the population it will serve. Hardware that delivers promised TFLOP/s is necessary but insufficient; the model running on that hardware must preserve the quality users depend on, and the data that shaped that model must represent the world it will encounter. Completing the validation stack requires turning from hardware to the model and data dimensions of the three-dimensional framework.

12.11 Model and Data Evaluation

A compressed model running on accelerated hardware can still fail if it was trained on biased data. System benchmarks can confirm that hardware delivers promised training throughput, inference latency, and power efficiency, but hardware validation alone cannot ensure deployment success. The optimization pipeline from Part III also included model compression (chapter 10) and data selection (chapter 9), each requiring its own validation. The remaining two dimensions of the framework address this gap: model benchmarks verify that compression preserved accuracy and critical model properties, while data benchmarks verify that training data enables robust generalization.

12.11.1 Model benchmarking

Model benchmarks validate whether compression techniques from chapter 10 preserved the properties that matter for deployment. This extends beyond top-line accuracy. A pruned model might maintain ImageNet accuracy while losing robustness to adversarial inputs. A quantized model might preserve average-case performance while degrading on rare but critical edge cases. A distilled model might match the teacher’s accuracy while losing calibration. Historically, benchmarks focused almost exclusively on accuracy, but compression makes multi-dimensional evaluation essential.

ImageNet links model benchmarking to the hardware story from figure 12.1: error rates fell as GPU-enabled architectures became practical. Figure 12.11 traces that progression from 28.2 percent error in 2010 to 3.57 percent in the ImageNet Large Scale Visual Recognition Challenge (Russakovsky et al. 2015). The introduction of AlexNet³⁵ reduced the error rate from 25.8 percent to 16.4 percent. Subsequent models like ZFNet, VGGNet, GoogLeNet, and ResNet³⁶ continued this trend, with ResNet achieving 3.57 percent (He et al. 2016a). This progression established the baselines against which model compression techniques are evaluated: a pruned ResNet must demonstrate how much accuracy it sacrifices for a given efficiency gain.




Figure 12.11: ImageNet Challenge Progression: AlexNet marks the largest one-year drop in the series; ZFNet, VGGNet, GoogLeNet, and ResNet extend that decline, defining the accuracy trajectory against which later compression trade-offs are judged. Sources: (Russakovsky et al. 2015; Krizhevsky et al. 2012; He et al. 2016a).

12.11.1.1 Accuracy metrics and their blind spots

The most common model metrics (accuracy, precision, recall, F1) each reveal different aspects of model behavior while hiding others, and understanding their blind spots is essential for compression validation. Top- k accuracy measures whether the correct label appears in the model's top- k predictions. Top-1 accuracy is strict; top-5 is lenient. The gap between them reveals model uncertainty: a model with 75 percent top-1 but 95 percent top-5 accuracy “knows” the answer is among a few candidates but struggles to commit. For deployment, the acceptable gap depends on whether downstream systems can use ranked predictions or require single answers.

Precision and recall matter when classes are imbalanced or errors have asymmetric costs (Sokolova and Lapalme 2009). A fraud detection model with 99 percent accuracy might have 10 percent recall on actual fraud (catching only one in 10 fraudulent transactions), a catastrophic failure despite high accuracy. Precision (of predicted positives, how many are correct?) and recall (of actual positives, how many were found?) expose these failures that accuracy hides.

Most insidiously, aggregate metrics hide subgroup failures. A model achieving 95 percent overall accuracy might achieve 60 percent on a critical demographic subgroup. The Gender Shades project (Buolamwini and Gebru 2018) revealed commercial gender-classification systems for facial analysis performing substantially worse on darker-skinned women than on lighter-skinned men, a disparity invisible to aggregate benchmarks. Disaggregated evaluation across deployment-relevant subgroups is essential; chapter 15 examines fairness evaluation systematically.

12.11.1.2 Calibration: When confidence scores matter

For many deployment scenarios, *how confident* the model is matters as much as *what* it predicts. A well-calibrated³⁷ model's confidence scores correspond to actual correctness probability: when it says “90 percent confident,” it should be correct 90 percent of the time.

Compression can shift calibration even when preserving accuracy, a critical concern when validating quantization techniques from section 10.4. A quantized model might maintain headline

³⁵ | **AlexNet:** The eight-layer CNN (60M parameters) that cut ImageNet top-5 error from 25.8 percent to 16.4 percent in 2012, trained on two GTX 580 GPUs with 3 GB memory each (Krizhevsky et al. 2012). AlexNet established a benchmarking paradigm that still informs vision evaluation: accuracy on a fixed dataset as the primary metric, with hardware configuration as a secondary specification. Later ImageNet results inherited this baseline comparison structure.

³⁶ | **ResNet (residual network):** Introduced by He et al. (2016a), skip connections enabled 152+ layer networks and achieved 3.57 percent top-5 ImageNet error (ensemble), surpassing the estimated human error rate reported in the ImageNet challenge context (Russakovsky et al. 2015). ResNet-50 became a common MLPerf Training reference workload because its moderate size (25.6M parameters) and well-understood compute profile (8.2 GFLOP per image) make it sensitive to both hardware and software optimizations without requiring multi-node setups (Mattson et al. 2020).

³⁷ | **Calibration:** From Arabic *qalib* (a mold for casting metal) via Latin *calibrare*, originally describing the adjustment of measuring instruments against known standards. In ML, calibration ensures predicted probabilities match empirical frequencies; Guo et al. (2017) formalize this concern for modern neural networks and show that temperature scaling is a simple effective post-hoc correction. The etymology is apt: just as an uncalibrated instrument produces precise but inaccurate measurements, an uncalibrated model produces confident but unreliable predictions, causing downstream systems that threshold on confidence scores to make systematically wrong decisions.

accuracy while becoming overconfident on examples it gets wrong. This matters because post-hoc calibration techniques such as temperature scaling can only correct the problem if calibration is measured explicitly (Guo et al. 2017).

Calibration failures create downstream problems. An overconfident model automates errors that it should defer (predicted 95 percent confidence but wrong 30 percent of the time). An underconfident model sends too many correct predictions to review instead of automating decisions it can handle (predicted 70 percent confidence but correct 95 percent of the time). **Expected calibration error (ECE)** measures the gap between confidence and accuracy across confidence bins; reliability diagrams visualize this correspondence.

12.11.1.3 Compression validation: The efficiency-quality frontier

Model compression (chapter 10) trades model capacity for efficiency. Validation must determine whether compression achieved an acceptable trade-off or damaged capabilities that matter.

Pareto frontier³⁸ evaluation determines whether a compressed model represents a good trade-off. Plotting accuracy against the target efficiency metric (latency, model size, energy) reveals the trade-off frontier. Models on the Pareto frontier *cannot* improve one metric without degrading the other; models below the frontier are dominated by better alternatives.

Different compression techniques make different efficiency-quality trade-offs. Quantization reduces precision while often preserving aggregate accuracy (Jacob et al. 2018). Pruning induces sparsity while retaining varying levels of test performance (Han et al. 2015; Gale et al. 2019). Distillation transfers knowledge from a larger teacher or ensemble into a smaller model (Hinton et al. 2015). Because these aggregate results do not establish preserved calibration or tail behavior, validation must measure those properties directly (Guo et al. 2017).

Calibration is a failure mode that aggregate accuracy can hide, and expected calibration error (ECE) is one common diagnostic. ECE compares predicted confidence with empirical accuracy, but it has no universal “good,” “moderate,” or “poor” thresholds. The estimate depends on the binning scheme, bin count, sample size, and prediction distribution. A benchmark must therefore freeze the estimator and compare it with a task-specific tolerance and uncompressed baseline. Compression can leave top-1 accuracy intact while materially changing ECE, which is why a compression protocol measures it directly.

Pareto optimality identifies nondominated choices; it does not establish that any choice clears the deployment’s absolute quality floor. Acceptable degradation depends on deployment context. A 2 percent accuracy drop might be acceptable for a recommendation system (users tolerate imperfect suggestions) but unacceptable for medical diagnosis (each error has significant consequences). Define accuracy thresholds before compression, then validate against them. The acceptance gate must cover aggregate accuracy, calibration, and deployment-critical slices. The MobileNetV2 lighthouse makes the complete INT8 validation protocol concrete.

³⁸ | **Pareto frontier:** Named after economist Pareto (1896), the frontier contains all solutions where improving one objective requires degrading another. In compression benchmarking, the frontier’s shape carries diagnostic information: a steep region means efficiency gains come cheaply (prune here), while a flat region means further compression costs disproportionate accuracy (stop here). Points below the frontier are strictly dominated and represent wasted capacity.



A dominated model loses on both axes of the Pareto frontier.



Lighthouse 12.3: MobileNetV2 INT8 compression

Returning to lighthouse 12.1, consider an illustrative validation protocol for INT8 quantization, grounded in MobileNetV2’s architecture (Sandler et al. 2018) and post-training quantization practice (Jacob et al. 2018). The values in table 12.18 are assumed:

Precompression baseline: MobileNetV2 achieves 71.8 percent top-1 accuracy on ImageNet at 3.5M parameters (14 MB FP32).

Notice in table 12.18 that aggregate accuracy barely changes after INT8 quantization to 3.5 MB, but calibration error and edge-case accuracy tell a different story. The INT8 model’s ECE rises from 0.031 to 0.089; whether that increase is acceptable depends on a task-specific tolerance under the fixed ECE protocol.

Table 12.18: MobileNetV2 INT8 Postcompression Metrics: Illustrative FP32 and INT8 values for top-1 and top-5 accuracy, calibration error, and edge-case accuracy, showing how aggregate accuracy can mask calibration and tail-case degradation.

Metric	FP32	INT8	Acceptable?
Top-1 accuracy	71.8%	70.9%	✓ (0.9 pp drop; below 1 percentage-point threshold)
Top-5 accuracy	91%	90.4%	✓
Calibration ECE	0.031	0.089	No (degraded)
Edge-case accuracy	68.2%	61.4%	No (drop of 6.8 pp)

Edge-case definition: Images with > 50 percent occlusion, < 100 lux lighting, or > 30° rotation from training distribution (approximately 5 percent of real-world inputs).

What this reveals: Under these assumptions, average-case accuracy looks acceptable (0.9 percentage-point drop), but calibration degrades and edge-case accuracy drops 6.8 percentage points. If the deployment uses confidence thresholds (for example, “only act if confidence > 85 percent”) or encounters many edge cases (unusual lighting, partial occlusions), INT8 MobileNetV2 could fail despite passing aggregate benchmarks.

Fix: Apply temperature scaling post-hoc to improve calibration (Guo et al. 2017). Temperature scaling learns a single scalar T_{cal} to divide logits before softmax: $\text{softmax}(z_i/T_{\text{cal}})$. In parallel, add edge-case examples to the test set to monitor that specific failure mode continuously.

The Lottery Ticket Hypothesis (section 10.3.1.7) provides concrete benchmarking data illustrating what Pareto-efficient compression looks like. Through iterative pruning, Frankle and Carbin (2019) found sparse subnetworks (“winning tickets”) in fully connected and convolutional networks that could match the original network’s test accuracy when trained in isolation.

The Lottery Ticket results reveal the shape of compression trade-offs: aggressive pruning can preserve accuracy for some architectures and tasks, but the acceptable sparsity point is empirical rather than universal. Compression validation should establish similar trade-off curves for each specific model and task, identifying where the model sits on the Pareto frontier and whether further compression yields meaningful efficiency gains or merely degrades quality.

12.11.1.4 Large language model benchmarks

The compression evaluation framework applies cleanly when the task has a stable label: classification accuracy, detection mAP, segmentation IoU. Large language models break that pattern. A team can choose a model because it scores well on a public benchmark, then discover in deployment that the model recognizes multiple-choice facts but cannot generate a grounded answer, responds too slowly for an interactive product, or produces confident unsafe text that the benchmark never stressed. LLM benchmarking therefore starts by naming the deployment failure that a score is meant to rule out.

The useful LLM metric taxonomy in table 12.19 is therefore a decision aid, not a leaderboard. Its rows use Massive Multitask Language Understanding (MMLU),³⁹ HELM (Holistic Evaluation of Language Models),⁴⁰ and perplexity⁴¹ as examples of scores that answer different deployment questions:

Table 12.19: Failure-Oriented LLM Benchmark Taxonomy: LLM metrics are useful when each score is tied to the deployment failure it is meant to rule out. Recognition, holistic behavior, corpus prediction, and responsiveness expose different risks, so no single score establishes production readiness.

Deployment failure to rule out	Metric or benchmark family	What the score reveals	What the score cannot prove
The model recognizes facts poorly	MMLU (Massive Multitask Language Understanding)	Broad factual and disciplinary knowledge across fifty-seven subjects, with scores interpretable against chance-level multiple-choice performance	Whether the model can generate grounded open-ended answers rather than choose among multiple-choice options

³⁹ | **MMLU (massive multitask language understanding):** Introduced by Hendrycks et al. (2020) with 15,908 multiple-choice questions across fifty-seven subjects. MMLU’s benchmarking limitation is its format: multiple-choice recognition is not the same task as open-ended generation, so an MMLU score should not be read as direct evidence that a model can produce grounded free-form answers in production.

⁴⁰ | **HELM (holistic evaluation of language models):** Stanford’s 2022 evaluation framework tested a broad set of models across seven dimensions (accuracy, calibration, robustness, fairness, bias, toxicity, efficiency) (Liang et al. 2022). HELM’s contribution is methodological: by evaluating models that score similarly on accuracy but diverge on calibration or toxicity, it demonstrates that single-metric leaderboards systematically hide failure modes that matter for production deployment.

⁴¹ | **Perplexity:** From Latin *perplexus* (entangled); in information theory, $2^{H(p)}$ where H is entropy. A perplexity of 10 means the model is “10-way confused” on average. The systems consequence is interpretive rather than direct memory accounting: perplexity measures held-out next-token prediction on a corpus, while serving memory pressure is governed by context length, batch size, model shape, and decoding state; KV-cache management is a separate serving problem (Kwon et al. 2023).

Deployment failure to rule out	Metric or benchmark family	What the score reveals	What the score cannot prove
The model is capable but unsafe	HELM (Holistic Evaluation of Language Models)	Accuracy alongside calibration, robustness, fairness, bias, toxicity, and efficiency	Whether one aggregate score captures the deployment risk; a model can be strong on accuracy and weak on calibration, safety, cost, or prompt stability
The model predicts its corpus well	Perplexity	Held-out next-token prediction on the same corpus; a perplexity of 10 means the model is “10-way confused” on average	Whether generated answers are helpful, safe, or grounded outside that corpus
The model feels slow in use	First-token latency, inter-token latency, and token throughput	Prompt-processing delay before generation starts and decode speed after generation begins	Whether a single throughput number hides poor interactive responsiveness, especially when batching improves throughput but worsens first-token latency

The responsiveness row deserves a concrete timing anchor because LLM benchmarks often report a single throughput number even though users experience generation in phases. A model can look efficient in tokens per second while still feeling slow if the first token arrives late, or it can improve first-token latency while producing the rest of the answer too slowly for an interactive workflow. The token throughput calculation turns those token-rate metrics into user-visible wall-clock time.

Token throughput turns that trade-off into wall-clock time. For a response of about 750 tokens, 25 tokens/s means 30 seconds of generation, while 100 tokens/s means 7.5 seconds. Time-to-first-token and inter-token latency must therefore be reported together: one captures responsiveness at the start of the exchange, and the other captures the rate at which the answer arrives.

The final failure is that a score may measure memory rather than capability. Benchmark contamination is a unique LLM risk because models trained on web-scale corpora may encounter benchmark questions during pretraining, inflating scores through memorization rather than skill (Xu et al. 2024). Leakage detection reframes this risk as something benchmark designers can test for rather than merely suspect. Temporal holdouts use content published after the training cutoff, dynamic benchmarks generate fresh instances continuously, and contamination tests ask whether the model recalls exact benchmark phrasing. These techniques keep the benchmark aligned with the deployment question instead of rewarding exposure to the test set.

12.11.2 Data benchmarking

Model benchmarks validate whether compression preserved model quality. Model quality, however, depends entirely on the data used to train and evaluate it, and this dependency creates the most insidious failure mode in ML deployment. A perfectly preserved model trained on biased or unrepresentative data will still fail in production. Data benchmarks validate whether the efficiency strategies from chapter 9 (active learning, curriculum design, data augmentation, and synthetic data generation) produced training sets that enable reliable deployment. This is often the last validation to fail and the hardest to diagnose: a model achieving excellent accuracy on held-out test data may collapse on production inputs that the training data never adequately represented.

Contemporary AI development has made clear that data quality can set a performance boundary that architecture changes alone cannot overcome. This recognition elevated data benchmarking from afterthought to critical discipline.

A data benchmark therefore starts with a protocol before it starts with a score. Define the deployment slice the model must serve, reserve a leakage-resistant holdout, verify duplicate and near-duplicate separation across partitions, set minimum coverage for rare classes and subgroups, audit label quality, and establish drift thresholds that determine when the benchmark no longer represents production. Only after those gates are explicit do aggregate metrics become interpretable.

Coverage metrics

The first question data benchmarking must answer is whether the training data represents the inputs the model will encounter. A model cannot learn patterns it has never seen, and the ways training data can fail to represent deployment reality are often subtle.

Consider class balance: a fraud detection dataset with 99 percent legitimate transactions and 1 percent fraud might produce a model that achieves 99 percent accuracy by simply labeling everything legitimate. The model is useless, but the accuracy metric looks excellent. Severe imbalance often requires mitigation through oversampling, class weighting, or threshold adjustment. More insidious is subgroup imbalance within classes: a dataset might have balanced positive and negative examples overall, but negative examples might be drawn predominantly from one demographic group, creating disparities invisible to aggregate class balance metrics.

Feature coverage presents an even harder challenge because it requires domain knowledge about what variations matter. A computer vision model trained exclusively on daytime images will fail on nighttime inputs; a natural language model trained on formal text will fail on colloquial language. Unlike class balance, which can be computed from labels alone, feature coverage requires understanding the deployment context. The lighting conditions the camera will encounter, the dialects users will speak, and the edge cases that exist in production but never appear in test sets all fall outside what labels alone can predict. These questions have no algorithmic answer; they demand collaboration between ML engineers and domain experts who understand the deployment environment.

For applications affecting people, demographic representation becomes a coverage dimension with ethical implications. Training data must represent the deployment population across relevant dimensions: age, gender, ethnicity, geography, language. A facial recognition system trained predominantly on one demographic group may underperform on others, even if aggregate accuracy metrics look acceptable. The challenge is that demographic metadata is often unavailable or unreliable, making representation gaps difficult to detect and measure.

Quality metrics

Even when training data covers the right inputs, the labels themselves may be unreliable. Studies consistently find 3–6 percent label error rates in major datasets, including ImageNet (Northcutt et al. 2021). These errors are not merely noise—they become learned ground truth. A model trained on data where wolves are occasionally labeled as dogs will learn the false rule that *some wolves are dogs*. The benchmark will report this as correct behavior because the model matches the (incorrect) labels.

For small datasets, manual audit of a random sample can estimate label accuracy. For large datasets, confident learning techniques identify likely mislabeled examples by finding cases where model predictions systematically disagree with labels. The intuition is that when a model confidently predicts a different label than the ground truth, either the model has learned something incorrect or the label is wrong. Detection, however, is only the first step; correction requires human review, and scaling human review to millions of examples presents its own challenges.

Inter-annotator agreement provides a different lens on label quality by measuring consistency across human labelers. Cohen's kappa or Fleiss' kappa quantify agreement beyond what chance would produce (Cohen 1960; Fleiss 1971). When agreement falls below conventional thresholds for tasks with clear ground truth, something is wrong: either the labeling guidelines are ambiguous, the task is inherently subjective, or labeler quality varies significantly. Landis and Koch's qualitative kappa bands are widely cited as a rough interpretive guide, though they should not replace domain judgment (Landis and Koch 1977).

The distinction between random and systematic errors matters enormously for their downstream effects. Unstructured label noise can be less damaging than a consistent bias because its errors do not all reinforce the same false rule, although modern models can still memorize it. Systematic errors (consistently mislabeling a particular subclass), in contrast, teach a coherent but wrong association. A dataset where all wolves photographed in snow are labeled "dogs" can produce a model that calls snowy wolves dogs, and adding more data with the same labeling rule only reinforces the error.

Distribution alignment

The final category of data benchmarking asks whether models will generalize from training conditions to deployment reality. This train-to-production alignment question is where the gap between benchmark performance and production performance most frequently emerges.

The standard assumption underlying held-out evaluation, that test data comes from the same distribution as training data, is routinely violated in practice. Test sets constructed years after

training data may reflect distribution drift as the world changes. Test sets from different geographic regions may reflect population shift. A model with strong held-out accuracy can drop sharply when deployed to a region or time period the test set did not represent. Once the i.i.d. (independent and identically distributed) assumption fails, held-out performance no longer identifies deployment performance and is often optimistic when the deployment distribution is harder.

The true test is train-to-production alignment, and this is far harder to measure because production data differs from training data in ways that held-out test sets often fail to capture. Production images come from different cameras with different characteristics. Production users come from different populations with different behaviors. Production inputs include edge cases that curated test sets systematically exclude. The WILDS⁴² benchmark (Koh et al. 2021) was designed specifically to evaluate models under realistic distribution shifts: hospital systems with different patient populations, wildlife cameras at different locations, satellite imagery from different time periods. On Camelyon17-WILDS, the reported ERM baseline achieved 93.2 percent in-distribution average accuracy and 70.3 percent out-of-distribution average accuracy.

Given these challenges, shift detection methods become essential for production monitoring. Statistical tests like the Kolmogorov-Smirnov test (Berger and Zhou 2014) or kernel-based two-sample tests such as maximum mean discrepancy (Gretton et al. 2012) can detect covariate shift—when the distribution of inputs changes even if the relationship between inputs and outputs remains stable. Monitoring model confidence distributions can detect when the model encounters inputs unlike anything in training. The goal is early detection: identifying distribution shift before it causes catastrophic performance degradation, enabling intervention through model updates, data collection, or deployment constraints.

Distribution alignment challenges highlight a persistent tension in ML development between two paradigms: *fixing the data and iterating on models*, or *fixing the model and iterating on data*. Figure 12.12 places these two paradigms side by side, revealing exactly where the feedback loop differs. In the model-centric diagram, the iteration cycle targets the architecture while the data remains static; in the data-centric diagram, the architecture stays fixed while the cycle targets data quality. The two approaches are complementary rather than a universal ranking of where improvement will come from.




Figure 12.12: Development Paradigms: Model-centric AI iterates on architecture while holding the dataset fixed, whereas data-centric AI iterates on annotations, diversity, and bias while holding the architecture fixed. The double arrow marks the approaches as complementary.

Data-centric AI reflects an important shift in understanding that challenges the “more data is always better” assumption: dataset composition matters alongside scale. Initiatives like DataPerf (Mazumder et al. 2023) and DataComp⁴³ have emerged to evaluate how dataset construction affects model performance systematically. In DataComp’s compute-controlled setting, a baseline that retained the top 30 percent of the candidate pool by CLIP-based filtering outperformed the unfiltered-pool baseline on aggregate downstream evaluation (Gadre et al. 2023). The result establishes the value of curation for that protocol, not a universal optimal fraction.

A persistent challenge in data benchmarking emerges from dataset saturation. When models achieve near-perfect accuracy on benchmarks like ImageNet, practitioners must distinguish whether performance gains represent genuine capability advances or merely optimization to existing test sets. As the timeline in figure 12.13 illustrates, widely tracked AI benchmarks have repeatedly crossed reported human baselines, making each corresponding benchmark less useful as a differentiator (Maslej et al. 2024).

⁴² | **WILDS:** Stanford’s 2021 benchmark of ten datasets with real-world distribution shifts: hospital changes (Camelyon17), wildlife camera location shifts (iWildCam), and satellite imagery temporal drift (PovertyMap). For the reported Camelyon17-WILDS ERM baseline, average accuracy fell from 93.2 percent in-distribution to 70.3 percent out-of-distribution, demonstrating that standard held-out evaluation can overestimate performance when the i.i.d. assumption fails.

⁴³ | **DataComp:** Introduced in 2023, DataComp fixes the model family, training code, and compute budget while participants vary dataset construction. Its filtering track isolates the effect of selecting examples from a common candidate pool, making curation strategies comparable without attributing every gain to additional training compute.




Figure 12.13: Dataset Saturation: The zero line denotes reported human performance; each curve's crossing marks a capability for which the static benchmark has less remaining discriminating headroom. The crossings do not distinguish genuine capability gains from optimization to the evaluation set. Sources: (Maslej et al. 2024; Kiela et al. 2021).

Dataset saturation and dynamic benchmarks

Figure 12.13 raises a critical methodological problem: when models surpass human performance on benchmarks, the result may reflect either genuine capability advances or optimization to static evaluation sets, and the two are difficult to distinguish from leaderboard scores alone. MNIST, introduced through the classic handwritten-digit recognition work of LeCun and colleagues (LeCun et al. 1998), illustrates the concern: static test images can contain dataset-specific artifacts that models learn to exploit. The question “Are we done with ImageNet?” (Beyer et al. 2020) generalizes this concern.

Dynamic benchmarking approaches like Dynabench⁴⁴ (Kiela et al. 2021) address saturation by continuously evolving test data based on model performance, ensuring that benchmarks remain challenging as capabilities improve. However, dynamic benchmarks complement rather than replace the coverage, quality, and distribution metrics described earlier: they prevent saturation but do not diagnose its causes.

12.11.3 Holistic system-model-data evaluation

Passing system, model, and data benchmarks independently is not enough. A system benchmark can validate hardware performance, a model benchmark can verify that compression preserved quality, and a data benchmark can assess training set representativeness, yet the deployed system can still fail because the three dimensions interact. Real-world AI performance emerges from that interaction, and optimizing one dimension can expose weaknesses in another.

Consider a concrete failure cascade: a team achieves excellent MLPerf Inference scores by deploying an INT8-quantized model on optimized hardware. System benchmarks pass. The quantized model, however, was validated only on ImageNet-distributed test data; deployment reveals accuracy degradation on factory-floor images with different lighting characteristics. Model quality benchmarks would have caught the quantization sensitivity. Further investigation shows the training data contained no images with industrial lighting—a data quality gap that no amount of system or model optimization can address.

This interdependence means that benchmark results from one dimension can be invalidated by failures in another:

- System success + Model failure: Hardware delivers promised throughput, but compression degraded accuracy below deployment thresholds
- System success + Data failure: Fast inference on representative inputs, but training data bias causes failures on demographic subgroups

⁴⁴ | **Dynabench:** Facebook AI Research's 2021 platform for dynamic benchmark generation, where humans craft adversarial inputs that fool current best models. Dynabench addresses the saturation problem, where very high accuracy on static benchmarks may reflect test-set familiarity rather than robust capability, but introduces its own trade-off: dynamic benchmarks are harder to compare across time because the evaluation set changes. Static and dynamic benchmarks serve complementary diagnostic roles.

- Model success + System failure: Accurate predictions, but latency variance under load violates SLA requirements
- Model success + Data failure: High accuracy on held-out test set, but distribution shift in production causes silent degradation

This interdependence echoes the D·A·M landscape introduced in chapter 1 (figure 1.4), but the benchmark lenses do not map one-to-one onto its axes. Data benchmarks primarily test the evidence supplied to learning; model benchmarks primarily test learned algorithmic behavior; and system benchmarks expose machine execution together with the data movement and software needed to drive it. Holistic evaluation verifies that assumptions made in one lens still hold where the D·A·M axes interact. The Part III optimization pipeline (data → model → hardware) creates dependencies that benchmarking must validate explicitly.

The D·A·M taxonomy provides a diagnostic framework for systematically identifying which axis limits performance. Section A.1 maps each axis to its binding physical constraint and the optimization pathway that relieves it, giving the first diagnostic step when a benchmark reveals underutilization. Table 12.20 formalizes this approach by crossing each D·A·M axis with the three fundamental bottleneck types; Appendix A gives the full diagnostic guide, including profiling utilities and efficiency screening indicators.

Table 12.20: D·A·M Bottleneck Diagnostic Matrix: Each cell describes a performance constraint symptom, and the row identifies which D·A·M axis to address. When performance stalls, this matrix turns the first diagnostic move into an axis-specific check of Data, Algorithm, or Machine.

Component	Compute-Bound	Memory-Bound	I/O-Bound
Data	Preprocessing too slow (augmentation, tokenization)	Dataset exceeds RAM (spills to disk)	Storage cannot feed GPU (disk throughput limit)
Algorithm	Model too large for hardware (FLOPs exceed capacity)	Activations exceed memory (batch size limited)	Gradient sync slower than compute (distributed training)
Machine	GPU utilization saturated (need faster accelerator)	Memory bandwidth saturated (need more HBM bandwidth)	Network/PCIe bandwidth saturated (need faster links)

The diagnostic power of this matrix becomes clear when benchmarks reveal unexpected results—particularly when performance falls short of expectations. If system benchmarks show low GPU utilization despite adequate hardware, the bottleneck likely lies elsewhere. For example, if only 30 percent GPU utilization is observed during training, an inefficient model architecture (Algorithm row) might initially be suspected, but profiling reveals that image augmentation runs on CPU and cannot keep up with GPU consumption (Data row, Compute-Bound column: “Preprocessing too slow”). Systematic diagnosis using this matrix prevents the common mistake of optimizing the wrong component.

Validation under controlled laboratory conditions differs from validation in production. In the laboratory, data distributions stay fixed, request patterns remain uniform, and systems run in isolation. In production, all three assumptions can break simultaneously—data drifts, traffic spikes unpredictably, and system components interact in ways that isolated benchmarks cannot capture. The final dimension of benchmarking asks whether laboratory results hold under operational conditions.

12.12 Production Considerations

A system that passes all three benchmark categories can still fail in production. The three-dimensional framework validated hardware performance, model quality, and data representativeness under controlled conditions—but production violates those conditions continuously. This gap between benchmark success and deployment success motivates a final benchmarking concern: validating systems under conditions that match operational reality.

12.12.1 From laboratory to production

Laboratory benchmarks establish what a system is capable of under ideal conditions. Production validation determines whether that system is performing correctly right now, under real conditions.

This distinction matters because laboratory benchmarks assume conditions that production systematically violates. Silent degradation poses the most insidious challenge: models can produce plausible but incorrect outputs without obvious error signals, and a recommendation system returning “reasonable” but suboptimal suggestions has no built-in error indicator. Dynamic workloads present a different failure mode: a system benchmarked at steady 1,000 QPS may fail when flash traffic events spike to 10,000 QPS, revealing that benchmark “throughput” assumed uniform request arrival rather than bursty production patterns. Data distribution shift compounds these problems over time, as production data evolves and diverges from training distributions—an image classifier trained on professional photos can degrade as users submit smartphone images with different lighting, angles, and compression artifacts. Finally, production imposes multi-objective constraints that benchmarks treat independently: accuracy, latency, cost, and resource utilization must all be satisfied simultaneously, and optimizing any one at the expense of others leads to deployment failure.

12.12.2 Bridging benchmark to deployment

Before deployment, validate benchmarking conclusions against production-representative conditions. Table 12.21 names the benchmark assumption, the production reality that violates it, and the validation step that closes the gap; the checkpoint that follows turns those rows into release-readiness actions.

Table 12.21: Predeployment Benchmark Checklist: Each row pairs an assumption baked into laboratory benchmarks with the production reality that violates it and the validation step that closes the gap. A system that passes every laboratory benchmark can still fail in production unless each row of this table has been independently verified against representative operational conditions.

Benchmark Assumption	Production Reality	Validation Approach
Uniform request arrival	Bursty traffic patterns	Load test with production trace replay
Clean, preprocessed inputs	Variable quality inputs	Evaluate on production data sample
Warm system state	Cold starts, cache misses	Measure cold-start performance
Isolated execution	Resource contention	Benchmark under realistic system load
Fixed model version	A/B testing, gradual rollout	Establish baseline for comparison

12.12.3 Production monitoring as continuous benchmarking

Production monitoring extends benchmarking from a one-time gate to a continuous process. The same principles apply (standardized metrics, reproducible measurement, statistical rigor) but the context shifts from “will this work?” to “is this working?”

Once a model is live, benchmarking becomes a rolling comparison against the baselines just established. The immediate checks stay concrete: whether the input distribution remains close to the benchmark distribution, whether latency and throughput stay inside the measured envelope, and whether model quality moves outside the expected range. Answering those checks requires the same measurement discipline as the offline benchmark, but now the measurements arrive continuously and under live traffic.



Checkpoint 12.4: Predeployment benchmark checklist

Before deploying a model based on benchmark results:

- ☐ Replay production traces: Use logged request patterns to validate throughput/latency under realistic conditions.
- ☐ Test with production data: Sample recent production inputs (respecting privacy) to verify accuracy holds.
- ☐ Stress test edge cases: Identify worst-case inputs and verify graceful degradation.
- ☐ Establish monitoring baselines: Document expected metric ranges for anomaly detection.
- ☐ Define rollback criteria: Specify quantitative thresholds for reverting to the previous model version.

The MLOps chapter later turns this measurement loop into release and recovery machinery: staged rollouts, shadow evaluation [running the new model beside production without serving its outputs], continuous validation, and rollback. At this point, the handoff is narrower. Benchmarking defines the baselines and failure thresholds; operations keeps measuring against them after deployment.

The same gap between benchmark conditions and production conditions explains why otherwise careful teams still make predictable mistakes. The final section names the misconceptions that turn benchmark success into deployment failure.

12.13 Fallacies and Pitfalls

Benchmarking creates false confidence when standardized measurement obscures deployment realities. Teams assume controlled evaluations predict production performance, but real systems face variability, resource constraints, and multi-objective trade-offs that benchmarks cannot capture, wasting engineering effort on systems optimized for evaluation rather than deployment.

Fallacy: *Benchmark performance directly translates to real-world application performance.*

The seductive clarity of benchmark rankings leads teams to select systems as though leaderboard position predicts production behavior. It rarely does. As section 12.3.1 demonstrates, ML systems exhibit inherent variability from data quality issues, distribution shifts, and resource constraints absent in controlled evaluation. In a representative failure scenario, a language model achieving 92 percent benchmark accuracy drops to 78–82 percent accuracy in production when processing user-generated text with spelling errors, informal language, and domain-specific terminology. An inference system with 15 ms mean latency on MLPerf experiences 150–200 ms p99 latency in production (10–13.3× degradation) due to concurrent load, garbage collection pauses, and network variability. Teams relying solely on benchmark rankings systematically underestimate deployment complexity, leading to failed launches and costly re-engineering.

Pitfall: *Optimizing exclusively for benchmark metrics without considering broader system requirements.*


Benchmark leaderboards incentivize aggressive optimization, but the optimizations that climb rankings often degrade the very characteristics production demands. This exemplifies Goodhart’s Law (section 12.10.4): when benchmark scores become optimization targets, they cease to be meaningful measures of system quality. In one illustrative scenario, a team reduces inference latency from 12 ms to 8 ms through aggressive quantization, improving MLPerf ranking by 15 positions while degrading calibration such that prediction confidence scores become unreliable for downstream decision-making. Another team improves ImageNet accuracy by 2.1 percent through extensive hyperparameter tuning but the optimized model consumes 40 percent more energy and exhibits 25 percent worse performance on out-of-distribution images from production cameras. Organizations rewarding benchmark rankings over deployment success systematically produce systems that excel in evaluation but fail in production.

Fallacy: *Single-metric evaluation provides sufficient insight into system performance.*

A single-number claim can seem seductively simple (“94 percent accurate” or “1,200 QPS fast”). But production success requires balancing multiple competing objectives that any single metric obscures. Modern inference systems demand evaluation across accuracy, latency, throughput, energy, and robustness dimensions (section 12.8.2). In an illustrative trade-off, a recommendation model achieving 94 percent accuracy with 180 ms p99 latency fails service-level objectives requiring p99 < 100 ms despite excellent accuracy. Conversely, a system optimized for 1,200 QPS throughput achieves this rate while consuming 4.2 W vs. 1.8 W for a slightly slower system at 1,000 QPS (2.3× power difference). For battery-powered edge devices, the 17 percent throughput loss enables 2.3× longer operation time. Different stakeholders prioritize different metrics: ML engineers focus on accuracy, infrastructure teams on throughput and cost, product managers on latency percentiles. Single-metric optimization systematically produces systems that excel on one dimension while failing deployment requirements on others.

Pitfall: *Using outdated benchmarks that no longer reflect deployment challenges and requirements.*

Benchmarks have inertia: teams continue reporting on established benchmarks after the results cease to provide useful discrimination. Saturation occurs when multiple approaches achieve near-identical



performance, eliminating useful comparison. ImageNet top-5 classification error decreased from 28.2 percent in 2010 (Russakovsky et al. 2015) to 3.57 percent by 2015 (He et al. 2016a), sharply compressing the headroom measured by that metric. As gaps narrow, the statistical confidence intervals discussed in section 12.10.1 and the result's deployment relevance matter more than leaderboard rank. Changing deployment contexts compound the problem: benchmarks designed for server hardware become misleading for edge devices with 10× less memory and 100× lower power budgets. Effective benchmarking requires retiring saturated benchmarks and developing evaluation frameworks matching target deployment realities.

Fallacy: *Research benchmarks predict production behavior under real traffic.*

Research benchmarks exist to compare algorithms under controlled conditions; production systems exist to serve users under variable ones. Treating the former as a production prediction often produces optimistic results because research benchmarks may omit resource contention, input-quality variation, and operational failure modes. Improper execution adds a separate source of error: omitting **warmup runs** can mix one-time initialization, just-in-time compilation, cache population, and memory allocation into steady-state latency; leaving cache state unspecified can make a nominal memory benchmark measure either warm-cache reuse or cold-memory access; and allowing dynamic voltage and frequency scaling (DVFS) to vary without reporting it makes small differences difficult to reproduce. Production systems face concurrent user loads, varying input quality, network latency, and system failures that degrade performance (section 12.10.2). A system achieving 800 QPS throughput in isolated benchmarks sustains only 400–500 QPS under production load with 90 percent utilization (37.5–50 percent degradation) due to queue contention and garbage collection pauses. Research benchmarks report model inference time (5–10 ms) while production end-to-end latency includes preprocessing, queuing, and postprocessing overhead totaling 50–100 ms. Production systems require 99.9 percent availability (43 minutes downtime per month) and graceful degradation under failures, characteristics research benchmarks often omit. Effective production evaluation requires operational metrics: sustained throughput under load, recovery time from failures, and complete latency breakdown.

Pitfall: *Using research benchmarks as production release gates.*

Teams sometimes promote a model because it passes the research benchmark, then discover only after launch that the benchmark never exercised the operational path. A release gate for a serving system must include load tests, tail-latency measurements, data-quality checks, failure drills, and rollback criteria. Research benchmarks remain useful for comparing algorithms, but production gates must measure the deployed system under the traffic, hardware, and failure conditions it will actually face.

12.14 Summary

Benchmarking completes Part III's optimization pipeline by validating whether the efficiency gains from data selection (chapter 9), model compression (chapter 10), and hardware acceleration (chapter 11) deliver in practice. The three benchmark lenses examine system execution, model behavior, and data representativeness separately, then test whether their assumptions survive when the complete system runs under deployment conditions.

The lenses reveal different failure modes. System benchmarks expose underdelivered throughput, tail latency, and thermal behavior. Model benchmarks test accuracy, calibration, robustness, and other properties that an optimization may alter. Data benchmarks examine coverage, label quality, leakage, and alignment with deployment. Standardized suites such as MLPerf Training and Inference provide comparable system evidence; model and data protocols supply the quality and representativeness evidence that a hardware result cannot establish.

Rigorous benchmarking is what distinguishes engineering claims from guesses. Practitioners who validate their optimizations rigorously, by measuring wall-clock latency rather than trusting FLOP counts, profiling tail latencies rather than averages, and testing on production-representative data rather than convenient benchmarks, build systems that perform as expected when deployed. As AI systems become increasingly influential in critical applications, this measurement rigor determines whether optimization claims translate into real-world impact.



Key Takeaways: Measuring what matters

- **Benchmarks validate co-design:** System, model, and data benchmarks expose hardware underdelivery, compression quality loss, and distribution mismatch. A system that passes only one axis can still fail when Data, Algorithm, and Machine constraints meet under production load.
- **Proxy numbers need boundaries:** Standardized run rules make comparisons honest, but fixed workloads are still proxies. Batch size, thermal state, input distribution, concurrency, and service-deadline windows decide whether a lab result survives the benchmark-production gap.
- **Granularity trades diagnosis for realism:** Micro-benchmarks isolate kernels, macro-benchmarks expose model-level costs, and end-to-end benchmarks capture user-visible behavior. Effective measurement stacks all three so teams can see both the symptom and the layer that caused it.
- **Tail latency is the benchmark:** Interactive systems fail at p95 and p99 before averages move. Reporting percentile latency under representative load prevents a benchmark from approving a system whose mean passes while its worst-served requests violate the SLO.
- **Amdahl caps every optimization claim:** A faster model cannot outrun the rest of the pipeline; if preprocessing is 50 percent of latency, an infinitely fast model yields only a $2\times$ system improvement. Benchmark the whole request path before celebrating kernel speedups.
- **Efficiency still needs quality evidence:** INT8 may cut raw weight storage $4\times$; the simplified component model estimates an energy reduction of about $6.6\times$, but whole-device measurement, calibration, subgroup robustness, and edge-case behavior decide whether the compressed model is deployable.

Each chapter in this part promised fewer FLOPs, a smaller model, or higher throughput. Benchmarking tests those promises against the complete system. The gap between a claimed improvement and a measured one reveals where Data, Algorithm, and Machine were assembled rather than matched. Amdahl's Law shows why an infinitely fast model can still leave the pipeline bounded by everything outside it, while tail latency shows why an average can pass even as the worst-served requests fail. This is co-design held to account. An ML system is engineered, not asserted, and only measurement on the real workload can tell the two apart.



What's Next: From lab to live

What can a system that passes every benchmark still fail to predict? Static benchmarks cannot fully capture traffic bursts, data drift, or cascading failures under live workloads. Part IV leaves the controlled benchmark environment for production, beginning with chapter 13.

IV

DEPLOYMENT AND OPERATIONS

Deployment Principles

The code can be correct and the benchmarks excellent, and yet the system may fail. Part IV moves from controlled environments to production, where ML systems face the additional threat of silent decay. A machine learning system can continue to produce outputs that are confident, well-formatted, and wrong as the world drifts away from its training distribution. At deployment, the data environment escapes the engineer's control and can stress the trained algorithm and the serving machine in ways no test set anticipated. Reliability is therefore a continuous control loop of D·A·M co-design rather than a one-time release gate. The principles here define the requirements and diagnostic models for that reliability.

III Principle 9: The verification gap

Invariant: Statistical verification holds only over a defined deployment population P_{deploy} , task distance d , tolerance τ , and target error ϵ , estimated by a finite-sample procedure at a stated confidence level:

$$\Pr_{(x,y) \sim P_{\text{deploy}}} [d(f(x), y) \leq \tau] \geq 1 - \epsilon$$

Implication: Deployment is not a one-way transfer; it is a control loop. Because no test suite can cover every possible real-world input, production systems must monitor their own uncertainty and fail gracefully when they drift outside their known performance envelope.

The verification gap means finite evaluation can only estimate behavior on a defined population with uncertainty. Those bounds may erode as production data diverges from the data used to set them.

III Principle 10: The statistical drift diagnostic

Invariant: Divergence between the production and baseline distributions does not by itself determine whether accuracy changes; over a measured deployment range, labeled outcomes can fit a local linear relationship:

$$\text{Accuracy}(t) \approx \text{Accuracy}_0 - \lambda \cdot \mathcal{D}(P_t \| P_0)$$

where P_0 is the baseline distribution, Accuracy_0 is the model's performance on P_0 , $\mathcal{D}(P_t \| P_0)$ is the statistical divergence between the current and baseline distributions, and λ is fitted from labeled outcomes rather than inferred from divergence alone. Consider a credit scoring model trained on 2020 borrower behavior. Two years later, inflation rises, interest rates change, and lending policies shift. The system still produces scores, but the statistical relationship between inputs and outcomes may have changed; labeled outcome monitoring determines whether real accuracy declined while conventional error logs remained quiet. Unlike many traditional software failures, which are often surfaced by crashes, exceptions, or explicit service-health signals, ML systems can fail silently because the *environment* changes even when the code and infrastructure remain unchanged. This first-order linearization applies only over the measured range; the relationship is model-dependent and may be nonlinear for large drift.

Implication: Observability must extend beyond system metrics (latency, errors) to combine statistical drift signals with labeled outcomes. A system can remain operational while prediction quality changes, but drift monitoring alone does not establish degradation.

External drift is not the only threat. Even when the world holds still, the serving pipeline itself can diverge from the model validated offline.



Principle 11: The training-serving skew diagnostic

Invariant: Training and serving code paths can diverge on identical inputs, and that divergence is measurable without labels. For scalar model scores, define training-serving divergence as the expected absolute difference between the two functions over a stated comparison-input distribution:

$$S_{\text{skew}} = \mathbb{E}[|f_{\text{serve}}(x) - f_{\text{train}}(x)|]$$

The exact relationship between S_{skew} and accuracy depends on the loss function, decision boundary geometry, and production distribution. Unexplained divergence weakens the assumption that offline validation estimates production behavior and can cause silent accuracy loss, but S_{skew} is not a universal accuracy-loss equation. This divergence arises from inconsistent preprocessing logic, different library implementations, stale feature values, or environmental state changes between the two code paths.

Implication: Feature consistency is an architectural requirement, not merely a best practice. Feature stores are more than caches; they can reduce skew by centralizing feature definitions and retrieval. Teams still need validation for freshness, point-in-time correctness, preprocessing, model-runtime, and postprocessing parity. Even subtle differences (PIL vs. OpenCV resize, FP64 vs. FP32 normalization) can compound to produce silent accuracy degradation that standard monitoring may not detect.

Beneath all these reliability concerns lies a nonnegotiable constraint: time. A medical imaging system that detects tumors with 99 percent accuracy but takes 30 seconds per scan may miss its workflow deadline. An autonomous vehicle perception model that classifies obstacles correctly but responds in 200 ms instead of 50 ms may miss its braking deadline. A correct result may have no operational value if it arrives too late. Latency-sensitive deployed models operate under a deadline, and exceeding it can be functionally equivalent to returning no prediction at all.



Principle 12: The latency budget principle

Invariant: Under serial request-path accounting, end-to-end latency is the sum of network, preprocessing, inference, postprocessing, and queueing delays, and a deployment meets its deadline only at a stated completion quantile:

$$L_{\text{lat,total}} = L_{\text{lat,net}} + L_{\text{lat,pre}} + L_{\text{lat,infer}} + L_{\text{lat,post}} + L_{\text{lat,queue}}, \quad \Pr[L_{\text{lat,total}} \leq D] \geq q$$

Implication: In latency-sensitive serving, the product selects the deadline and the acceptable completion quantile, such as 0.95, 0.99, or 0.999, and throughput is optimized within that constraint. Serving systems must implement tail-latency controls (for example, bounded-delay dynamic batching or hedged requests). Serving systems must be willing to sacrifice overall throughput to meet the latency deadline of the oldest request in the queue.

A system can satisfy its latency SLOs, detect measured distributional shifts, and maintain training-serving consistency while still causing systematic harm. The previous principles address silent failures in correctness and service quality; this one addresses failures that degrade *equity* through silent amplification.



Principle 13: The bias feedback model

Invariant: When outputs influence future inputs, today's disparity becomes an input to the next training cycle, so disparity evolves as a dynamical system rather than staying fixed. Let $\Delta_g(k) \geq 0$ denote the disparity magnitude for group g relative to a specified reference group after k fixed feedback intervals. With no effective intervention and an approximately constant dimensionless feedback factor $\alpha_{fb} > 0$, disparity may evolve as:

$$\Delta_g(k) \approx \Delta_g(0) \cdot \alpha_{fb}^k$$


where $\Delta_g(0)$ is the initial disparity magnitude relative to that reference group and α_{fb} is fitted per feedback interval. Consider a loan approval model that denies credit at higher rates to applicants from historically underserved communities. Denied applicants cannot build credit history, which makes future applications weaker, which increases future denial rates. The model's accuracy on its training distribution remains stable, but the population it serves has been reshaped by its own decisions. When $\alpha_{fb} > 1$ under these assumptions, the feedback loop is self-reinforcing; when $\alpha_{fb} \leq 1$, the modeled dynamics are stable or damped. Real deployments may also be nonlinear, saturating, or changed by intervention.

Implication: Fairness is not a postdeployment audit; it is a **stability constraint** on the deployment control loop. Where consequential impact warrants it and collecting or using group attributes is lawful and justified, systems must monitor **disaggregated performance metrics** across relevant demographic groups with the same rigor applied to latency percentiles, because a bias regression can be invisible to aggregate accuracy just as a tail-latency violation can be invisible to mean latency.

Together, the five principles define a deployment control loop that establishes what the evidence supports, observes the live system and environment, acts only when a measured failure warrants intervention, and re-evaluates the relevant service, predictive, or subgroup criterion. They separate diagnosis from response. A signal narrows the investigation but neither prescribes an automatic intervention nor proves that quality has already degraded.

Part IV makes that control loop concrete. Model Serving builds infrastructure that meets latency budgets under variable traffic. ML Operations detects drift and skew while preserving the evidence needed for diagnosis and response. Responsible Engineering treats fairness and feedback as measurable deployment dynamics rather than a one-time compliance check. The volume closes by connecting these production controls to the quantitative bounds and models developed in the preceding parts.

Model Serving




- 13.1 Serving Paradigm
- 13.2 Serving Load, Latency, and Architecture
- 13.3 Serving System Architecture
- 13.4 Request Lifecycle
- 13.5 Queuing Theory
- 13.6 Model Lifecycle Management
- 13.7 Throughput Optimization
- 13.8 LLM Serving
- 13.9 Inference Runtime Selection
- 13.10 Node-Level Optimization
- 13.11 Economics and Planning
- 13.12 Fallacies and Pitfalls
- 13.13 Summary

Purpose

Why does serving change the optimization priorities that made training successful?

Training and serving emphasize different operating objectives. Training usually maximizes throughput across large batches and long runs, while serving must answer individual requests within latency targets. Training amortizes hardware costs across many examples; serving pays a tax on each request, where small inefficiencies compound into operational debt. This shift is why models that train efficiently may still serve poorly. The live system has an end-to-end latency budget, so queuing, preprocessing, network transit, and postprocessing can erase an optimization confined to the model. Batch-heavy architectures and memory-intensive optimizations designed to saturate accelerators may not suit bursty, latency-critical, cost-sensitive production traffic. Serving, however, is more than a latency problem. A serving system must handle traffic that varies between peak and trough, reserve headroom for failures, shed load or degrade gracefully under overload, and remain predictable as request mix and demand change over time. Models that proved their value during training and survived compression and benchmarking eventually arrive at the serving layer, the deployment and integration stage of the ML lifecycle. The question shifts from whether a model works to whether it works reliably at scale under production conditions. Serving infrastructure is where ML systems meet users, and sustaining that interaction requires different engineering from creating the model. It is also where the trained algorithm meets live data within the machine's latency budget, bringing all three D-A-M constraints together on each request.





Learning Objectives

- Explain serving inversion from training throughput to per-request latency, headroom, and tail behavior
- Decompose request latency across serialization, preprocessing, inference, queuing, post-processing, and network overhead
- Apply queueing laws and simple queue models to plan capacity against percentile latency targets
- Diagnose training-serving skew and cold starts from mismatched preprocessing, model loading, or cache behavior
- Select batching, load shedding, autoscaling, and runtime strategies for traffic patterns and latency budgets
- Evaluate LLM serving bottlenecks using token latency, KV-cache memory, and continuous batching constraints
- Calculate cost per inference from precision, hardware utilization, replica count, and runtime throughput

13.1 Serving Paradigm

Serving begins where benchmarking stops: a model that performed under controlled measurement must now produce predictions under an operating workload. Cloud, Edge, Mobile, and TinyML each impose distinct serving challenges. In online serving, the priority often inverts from throughput over a long run to latency for each request. This serving inversion has concrete engineering implications that ripple through the whole stack. In the iron law of ML systems, the latency term (L_{lat}) now includes request scheduling, network round-trips, and system orchestration and may become the binding constraint. Controlled benchmarks establish performance under known conditions; serving faces traffic patterns no benchmark can fully anticipate. Quantization can reduce model size; serving must confirm that such optimizations preserve accuracy under real traffic distributions. Together these revalidations change the priorities of data, algorithm, and machine once requests arrive under an operational latency budget.

THE D-A-M taxonomy makes the inversion visible. For an online service, the data constraint shifts from training-set volume to the rate, freshness, and shape of live requests. The algorithm is ordinarily fixed for a deployed model version rather than updated through backpropagation on the request path. The machine constraint shifts from maximum utilization to **headroom**: operating below saturation leaves capacity for traffic spikes, while a saturated accelerator turns small load changes into tail-latency failures. The required margin depends on the arrival process, service-time distribution, and service-level objective (SLO). Serving therefore optimizes useful completed work under a latency promise rather than fully occupied hardware.

That promise ties the remaining parts of the serving stack together. Request routing, preprocessing, model execution, postprocessing, batching, caching, runtime selection, and capacity planning all compete for the same latency budget. The central engineering task is to decide which work belongs in the live request path, which work can move outside it, and how much headroom the system must reserve before useful throughput becomes fragile.

13.2 Serving Load, Latency, and Architecture



Example 13.1: The 'Black Friday' traffic spike

Scenario: An e-commerce recommendation serving system experiences a 10× traffic surge (10,000 QPS) during Black Friday sales events.

Diagnosis: When server utilization approaches 100 percent, queue lengths explode nonlinearly, increasing latency to 10 s and causing client request timeouts.

Systems lesson: High average throughput does not prevent queueing collapse under traffic surges. Depending on the service, preserving a tail-latency objective may require load shedding, graceful degradation, or autoscaling before utilization reaches the workload’s queueing knee.

Online serving systems may receive variable, unbatched request streams while still needing predictable response times across diverse physical environments. When a traffic spike exceeds a system’s reserved headroom, performance degrades nonlinearly.

Figure 13.1 shows that latency remains manageable at moderate utilization and then rises rapidly as the system approaches saturation; this is why latency-sensitive systems reserve headroom rather than planning for a permanently saturated accelerator. Section B.2.1 gives a mathematical treatment of long-tailed distributions and explains why high-percentile latency matters when an SLO covers the slowest requests. The curve is a simple queueing approximation intended for intuition rather than a specific workload.




Figure 13.1: The Tail Latency Explosion: Request latency vs. serving utilization ρ_{serv} . Mean latency (blue) and the p99 approximation (red) both grow nonlinearly as utilization approaches saturation; the 70 percent marker is an illustrative planning knee rather than a model-derived breakpoint. This uses the simple M/M/1 approximation introduced later in section 13.5 (p99 $\approx 4.6 \times$ mean), so the curve is illustrative rather than workload-specific.

Beyond the technical limits of latency, serving economics have also changed rapidly. As models become more efficient and hardware becomes more specialized, the cost per inference has fallen.¹ Facebook’s experience at fleet scale illustrates the magnitude of this serving cost problem.



Example 13.2: The inference tax at Facebook

Scenario: Production ML serving systems at Meta execute tens of trillions of inferences daily across global user-facing services (Hazelwood et al. 2018).

Diagnosis: Recurring inference serving compute and energy costs dwarf initial offline model training expenses. Live request arrival patterns and strict latency bounds constrain GPU batching.

Systems lesson: Model architectures cheap to train can prove unviable to serve at fleet scale. Production serving infrastructure requires hardware-software co-design to balance tail latency, batch size, and recurring operational total cost of ownership.

The same serving-economics pressure appears in public API prices. The log-scale price trajectory in figure 13.2 captures the speed of this cost collapse by tracking representative public API list-price snapshots as a market proxy. Vendor prices change frequently, so these points should be read as historical provenance for the trend rather than as current purchasing guidance (OpenAI 2023b,

¹ | **Jevons paradox:** William Stanley Jevons observed in 1865 that efficiency improvements in coal-powered steam engines increased total coal consumption by making steam power economically viable for applications previously too costly (Jevons 1865); the same dynamic can apply to AI inference: each 10 $\times$ cost reduction opens application classes that were economically infeasible at the previous price point, expanding aggregate demand by more than the efficiency gain. This is why cheaper inference can increase, not decrease, total GPU fleet demand. Efficiency and demand are often complements in AI, not substitutes.

2023a, 2024; OpenAI Developer Community 2024; Anthropic 2024; Google Developers Blog 2024; DeepSeek 2024). Each order-of-magnitude drop changes which applications are feasible.




Figure 13.2: Intelligence Deflation: Representative public API input-token list prices per 1M tokens (\$) over time (Log Scale). Eight model-version points are snapshots collected from the public pricing pages of OpenAI, Anthropic, Google, and DeepSeek between 2023 and 2025; the Legacy GPT-3.5 point is an illustrative baseline rather than a documented model-version snapshot. The series is intended as a market trend indicator, not a controlled or current-price comparison. API token-processing prices have collapsed by multiple orders of magnitude, transforming the economics of automated AI workflows.

Two pressures now frame the serving problem: tail latency that explodes once utilization passes the queueing knee, and per-inference economics that fall by orders of magnitude as efficiency improves. Together they force a formal definition of serving built around latency rather than throughput.

 Definition 13.1: Model serving

Model serving is the operational phase that provides model predictions to end-users or downstream systems under latency, throughput, availability, and cost constraints.

1. **Significance:** Online serving turns training’s throughput priority (η_{hw}) into a request-latency constraint (L_{lat}), requiring a stack designed around the percentile named by its SLO.
2. **Distinction:** Unlike model training, which commonly processes scheduled batches, serving may handle stochastic online requests, synchronized streams, or offline batches.
3. **Common pitfall:** A frequent misconception is that serving is “just the forward pass.” In reality, it is a systems problem: model execution is only one component of a stack that may include request routing, load balancing, and data transformation.

The SLO² defines the latency target that shapes every architectural decision in the serving stack, including how the system budgets time across preprocessing, model execution, postprocessing, and transport. When an inference server receives a request, model execution represents only one segment of an end-to-end processing pipeline spanning network transport, serialization, CPU pre/postprocessing, and queue management. Figure 13.3 shows that raw inputs pass through preprocessing (traditional computing), neural network inference (deep learning), and postprocessing (traditional computing) before producing final outputs. Any of these stages can become the latency bottleneck. Section 13.4.1 quantifies exactly where time goes, revealing a counterintuitive result about which stages dominate.

The pipeline turns serving into an orchestration problem: preprocessing, model execution, postprocessing, and transport all compete for the same latency budget. Before optimizing any one stage, the system must decide whether predictions are computed ahead of time or on demand.

2 | **Service level objective (SLO) vs. service level agreement (SLA):** An SLO is an *internal* target (for example, “p99 latency under 50 ms”); an SLA is an *external* contractual commitment whose violation may carry remedies or penalties. Teams often set SLOs tighter than SLAs to preserve a safety margin. For ML serving, task quality and inference latency can both contribute to SLOs, creating multidimensional targets where improving one dimension (for example, deploying a larger model for accuracy) can violate another (latency).




Figure 13.3: The Inference Pipeline: ML serving systems transform raw inputs into final outputs through sequential stages: preprocessing, neural network computation, and postprocessing. The neural network represents just one component; preprocessing and postprocessing make end-to-end latency depend on more than neural-network execution.

13.2.1 Static vs. dynamic inference

Before optimizing *how* to reduce inference latency, the system must decide *when* predictions are computed. An early architectural decision is whether predictions happen before or during user requests (Google 2024b). This choice shapes system design, cost structure, and capability boundaries.

13.2.1.1 Static inference

Static inference (also called offline or batch inference) precomputes predictions for anticipated inputs and stores them for retrieval. Consider a recommendation system that generates predictions for all user-item pairs nightly. When a user requests recommendations, the system retrieves precomputed results from a lookup table rather than running inference. This approach moves compute out of the request path, enables offline quality checks, and can reduce serving costs for predictable inputs. However, static inference needs either a fallback online path or a refreshed batch computation when requests include unanticipated inputs or newly updated models.

13.2.1.2 Dynamic inference

Dynamic inference (also called online or real-time inference) computes predictions on demand when requests arrive. This can handle previously unseen inputs within the service’s supported schema and reflects a new model version once that version is deployed. The cost is a live latency budget that constrains model complexity and demands robust monitoring infrastructure.

Example 13.3: The cost of latency

Latency constraints directly dictate infrastructure costs. Consider a GPU server renting for \$4/hour.

Scenario A (low latency): Batch size 1.

- Latency: 5 ms.
- Throughput: 200 req/s.
- Cost per million queries: \$5.56.

Scenario B (high throughput): Batch size 8.

- Latency: 10 ms (doubled due to batching overhead).
- Throughput: 800 req/s (quadrupled due to parallel efficiency).
- Cost per million queries: \$1.39.

Systems insight: Reducing latency from 10 ms to 5 ms increases the hardware bill by 300 percent. Engineers must quantify whether that speedup generates enough business value to justify the 4× cost increase.

For our ResNet-50 image classifier, consider two deployment scenarios. A static approach suits a photo organization app that preclassifies all images in a user’s library overnight. With 10,000 photos and 5 ms inference each, batch processing takes ~50 s total, and users see instant classification when browsing. A dynamic approach suits a content moderation API that must classify user-uploaded

images in real-time, with each image requiring the full preprocessing→inference→postprocessing pipeline and a 100 ms latency budget. A production image-classification system can combine the approaches: frequently requested images, such as popular products or known memes, are preclassified and cached, while novel uploads trigger dynamic inference.

The choice between static and dynamic serving has direct economic implications. A stricter latency requirement can raise infrastructure costs when it requires smaller batches, additional replicas, or more capable hardware. Quantifying this trade-off in dollar terms reveals whether the latency reduction justifies that premium.

Many production systems combine both approaches. Common queries can hit a cache populated by batch inference while uncommon requests trigger dynamic computation. Understanding this spectrum matters because it determines which subsequent optimization strategies apply. Static inference optimizes for throughput during batch computation and storage efficiency for serving. Dynamic inference optimizes for per-request latency under concurrent load, which requires understanding *where* time goes within each request.

The static-vs.-dynamic decision is the first of several architectural choices that shape serving system design. Equally important is *where* the model executes, since deployment context constrains every subsequent optimization.

All of the cost analysis in this section assumes a traditional forward pass: a fixed computation graph that executes once per request and produces a result. A new class of models upends that assumption by deliberately increasing the amount of computation spent per query, trading latency for answer quality, and the serving cost implications are substantial.

Systems Perspective 13.1: Looking ahead: Spending more compute per query

Traditional serving optimizes for minimizing latency ($L_{\text{lat}} \rightarrow 0$). Here, the optimization target expands to include answer quality. Some inference-time-compute systems deliberately spend more compute cycles to improve answer quality. Token generation is often memory-bandwidth bound at small batches, but larger batches or long contexts can shift the bottleneck toward compute or KV-cache traffic. These systems may generate far more tokens per request, including intermediate reasoning or search tokens, increasing the total compute and energy spent per query.

Whether a system spends one forward pass or many reasoning steps per query, deployment context still determines the feasible latency and cost envelope. That context is the next variable.

13.2.2 The spectrum of serving architectures

Although “serving” often implies a networked server processing API requests, the architectural pattern varies by deployment environment. Section 2.1 introduced the four deployment paradigms (Cloud, Edge, Mobile, and TinyML) and the physical constraints (the light barrier, the power wall, and the memory wall) that give rise to them. Architecturally, serving spans a continuum from **centralized networked microservices** (which pool accelerators to optimize throughput and cost-per-query at the expense of network round-trip overhead and cold-start latency) to **application-embedded edge serving** (which run models inside client processes via direct function calls, eliminating network latency and preserving privacy at the expense of strict device memory and power constraints). Those constraints do not disappear at serving time; serving adds latency SLOs and cost pressure on top of hardware limits that training could absorb through patience. The same model may require different serving strategies depending on where it executes.

13.2.2.1 Networked serving (cloud/data center)

In networked serving, the model runs as a standalone service (microservice), matching the cloud deployment paradigm that section 2.5 characterized as trading latency for larger pooled compute. The primary interface is the network through protocols such as HTTP or gRPC, so network transfer and serialization join model execution as possible binding constraints. Data-center hardware such as NVIDIA GPUs (V100, A100, H100), Google Tensor Processing Units (TPUs), and AWS Inferentia supports high-throughput batching and concurrency, but cold start can still stretch from seconds to

minutes because container startup, model loading, and warmup sit outside the steady-state inference path.

13.2.2.2 Application-embedded serving (mobile/edge)

In application-embedded serving, the model runs within the user application process (for example, a smartphone app using CoreML or TensorFlow Lite), following the embedded paradigm that section 2.6 and section 2.7 analyzed for its latency, privacy, and offline advantages. There is no “server.” The interface is a function call, so optimization focuses on energy and responsiveness (SingleStream) rather than shared-server throughput.

An important advantage is **zero-copy inference**: when data moves through a system, each copy consumes CPU cycles and memory bandwidth. In cloud serving, a camera frame might pass from a network buffer to application memory, a preprocessing buffer, GPU-accessible host memory, and finally accelerator memory. A mobile pipeline can remove some of these copies when its camera, image signal processor, CPU, and NPU exchange compatible shared buffers. This reduces latency and energy, although preprocessing or format conversion may still involve the CPU or image signal processor. The mechanism requires coordinated hardware and software support, as in unified-memory system-on-chip designs such as Apple A- and M-series processors and Qualcomm Snapdragon.

Typical hardware includes mobile NPUs (Apple Neural Engine, Qualcomm Hexagon) and embedded GPUs (Jetson). When the model is already resident and compiled, startup can be measured in milliseconds; first inference may take much longer if it triggers loading or just-in-time compilation. Mobile power and thermal envelopes are device-specific, and prolonged inference can trigger throttling.

13.2.2.3 Bare-metal serving (TinyML)

In TinyML serving, the model is commonly compiled into microcontroller firmware, reaching the extreme end of the deployment spectrum that section 2.8 introduced as ubiquitous sensing at microwatt power budgets. Many deployments have no general-purpose operating system or runtime allocator; “serving” is a tight loop that reads sensors and invokes an interpreter. Optimization focuses on fitting weights and preallocated working memory, such as a Tensor Arena, in static random-access memory (SRAM), and request batching is generally infeasible. Typical hardware includes ARM Cortex-M series, ESP32, and specialized TinyML accelerators. Startup can be extremely short when weights remain in flash and working memory is preallocated, while the power budget ranges from microwatts to milliwatts for battery operation over months or years.

The deployment tiers differ first in their operating envelopes. Table 13.1 compares the latency, batch, memory, power, update, failure, and monitoring constraints that follow from each environment.

Table 13.1: Serving Architecture Spectrum: Representative deployment envelopes for cloud, mobile, and TinyML serving. The latency, batch-size, memory, and power ranges are design guides rather than universal limits; update, failure, and monitoring rows summarize the operating pattern associated with each environment.

Characteristic	Cloud/Data center	Mobile/Edge	TinyML
Latency Target	10–100 ms	20–50 ms	1–100 ms
Batch Size	1–128 (dynamic)	1 (fixed)	1 (fixed)
Memory	16–80 GB VRAM	2–8 GB shared	256 KB–2 MB SRAM
Power	300–700 W	1–10 W	1–100 mW
Update Mechanism	Container deploy	App store update	Firmware over-the-air (OTA)
Failure Mode	Retry/failover	Graceful degradation	Silent or reset
Monitoring	Full telemetry	Limited analytics	Heartbeat only

Systems Perspective 13.2: ResNet-50 across the serving spectrum

Systems insight: The “same model” claim is misleading: the cloud and mobile tiers share the ResNet-50 graph but use different precision and runtimes, while the TinyML tier cannot run

ResNet-50 and substitutes an architecture designed for its constraints. Treating these as one model hides the work that makes each deployment possible.

Applying those envelopes to one image-classification workload, table 13.2 shows why cloud and mobile can adapt ResNet-50 while TinyML must substitute MobileNetV2.

Table 13.2: Image Classification Across the Serving Spectrum: Illustrative side-by-side comparison of cloud, mobile, and TinyML serving for the same image-classification task, showing how model format, latency, throughput, memory footprint, and energy budget differ across deployment contexts.

Dimension	Cloud	Mobile	TinyML
Model format	TensorRT FP16 engine (51.2 MB)	TensorFlow Lite INT8 (25.6 MB)	Not feasible (25.6 MB); alternative: MobileNetV2 INT8 (3.5 MB)
Inference (batch-1)	1.4 ms (batch-16: 14 ms)	12 ms (NPU), 45 ms (CPU)	120 ms
Throughput	1,143 img/s (batched)	~80 img/s (single-stream)	~8 img/s
Memory	2 GB VRAM (batch-32)	150 MB peak (shared with app)	320 KB arena (fits in 512 KB SRAM)
Energy/inference	—	0.8 mJ (NPU), 4.2 mJ (CPU)	12 mJ

13.2.3 The load balancer layer

When traffic exceeds what a single machine can handle, cloud and data center deployments that run multiple replicas of the same model require an additional infrastructure layer: the load balancer. Production serving systems place load balancers between clients and model servers, providing three essential functions for serving infrastructure.

Request distribution, the first function, routes incoming requests to available model replicas using algorithms like round-robin or least-connections. For latency-sensitive ML serving, algorithms that route away from slow or overloaded replicas improve tail latency. The second, health monitoring, continuously verifies that replicas are ready to serve, routing traffic away from unhealthy instances. For ML systems, health checks must verify both process liveness and model readiness, confirming that weights are loaded and warmup is complete. The third, deployment support, enables safe model updates by gradually shifting traffic between versions instead of treating release as an all-at-once switch. Section 14.5.1.1 later turns that basic traffic-shift idea into full deployment and validation strategies.


For single-machine serving with multiple model instances, such as running several Open Neural Network Exchange (ONNX) Runtime sessions, the framework and operating system handle request queuing. The full complexity of load balancing becomes necessary when scaling to distributed inference systems, where multiple machines serve the same model. The implementation details of request distribution algorithms and multi-replica architectures belong to that distributed context.

When capacity planning considers “the server” in this single-machine serving analysis, it means the machine’s model serving capacity. The queuing dynamics analyzed in section 13.5 apply to understanding single-machine behavior and determining when scaling to multiple machines becomes necessary.

While load balancers distribute requests *across replicas*, achieving predictable latency also requires controlling what happens *within each machine*. The operating system environment introduces its own sources of variability.

13.2.4 Deterministic latency and resource isolation

An inference server does not operate in isolation. On a single machine, the operating system manages multiple competing processes (logging agents, monitoring tools, and system interrupts) that can intermittently steal CPU cycles from the inference pipeline. These “noisy neighbors” are a primary source of **latency jitter**, where the time required to process identical requests varies significantly, causing the 99th percentile (p99) latency to spike even when the hardware is underused. Resource



contention can therefore widen tail latency through service-time jitter, a different mechanism from the utilization-driven queue growth in figure 13.1.

Achieving deterministic performance on a single node requires reducing interference from the operating system's normal resource-sharing behavior. Predictable serving systems such as Clockwork show that deep neural network (DNN) inference can meet tight request-level SLOs when scheduling and execution are controlled carefully (Gujarati et al. 2020). CPU affinity (pinning) is one local isolation tool: it restricts the inference server's threads to specific physical cores so latency-sensitive work is less exposed to thread migration and cache-locality loss. Pinning can reduce one source of latency jitter, but it is part of a broader resource-isolation strategy rather than a complete solution.

Memory locking (`mlock`) addresses a related but distinct source of jitter. By default, the OS can page eligible host memory regions to disk under memory pressure. If the CPU or GPU's direct memory access (DMA) engine accesses a region that has been paged out, the transfer stalls until the data is faulted back into RAM. Locking host-resident weights, staging buffers, or an offloaded KV cache prevents these stalls, though locked pages cannot be reclaimed by other processes. `mlock` does not lock accelerator high-bandwidth memory (HBM).

The third technique, interrupt shielding, completes the isolation picture. Network and storage interrupts routed to inference cores can preempt GPU command submission at unpredictable moments. Steering these interrupts to noninference cores ensures that bursts of incoming traffic do not disrupt the GPU's command stream, which is particularly important for maintaining stable tail latency under load.

These isolation principles transform a simple “model script” into a deterministic service, a transition essential for safety-critical applications like autonomous driving or real-time industrial control. The deployment spectrum, load balancing, and resource isolation define *where* models serve and *what* infrastructure supports them. The remaining question is *how* the serving software itself is organized, specifically what components comprise an inference server and how they coordinate to turn irregular user traffic into efficient hardware utilization.

13.3 Serving System Architecture

User requests arrive in unpredictable bursts, one millisecond apart, then five seconds of silence, while accelerators demand steady, uniformly-sized batches. Bridging this gap requires more than a Python script calling `model.predict()`; it requires a specialized software architecture that absorbs traffic variability, forms efficient batches, and keeps hardware saturated without violating latency SLOs.

13.3.1 Internal architecture and request flow

Model optimization focuses on the mathematical artifact, while model serving requires a specialized software architecture to manage high-frequency request streams and hardware utilization. An inference server³ (such as NVIDIA Triton or TensorFlow Serving) is not a simple wrapper around a model script; it is a high-performance scheduler that manages concurrency, memory, and data movement.

The internal anatomy of these servers reveals how they bridge the gap between irregular user traffic and the highly regular, batch-oriented requirements of accelerators. Every request traverses a multi-stage pipeline designed to maximize hardware throughput while minimizing latency overhead. Figure 13.4 separates the six stages so each component's role in absorbing traffic, queueing, batching, and accelerator execution is explicit.

This architecture serves three functions. First, concurrency management: servers use asynchronous event loops or thread pools to handle many concurrent client connections without blocking, so network I/O need not idle the accelerator. Second, request transformation: the server converts network payloads, such as JSON or Protobuf, into the tensor formats required by the selected runtime. Image tensors, for example, can use NCHW⁴ (batch, channels, height, width) or NHWC (batch, height, width, channels). The faster layout depends on the operator kernels, data type, hardware, and runtime; an unnecessary conversion between layouts can itself become part of the latency budget.

Third, model management: inference servers manage the lifecycle of loaded model artifacts, including loading weights into VRAM, tracking which artifact version is active, and completing

³ | **Inference server:** Google's TensorFlow Serving (Olston et al. 2017) helped establish the separation of model logic from serving infrastructure; NVIDIA's Triton (NVIDIA 2024d) extends this pattern across multiple model frameworks and backends. The critical design insight is that a scheduler and dynamic batcher turn irregular single-request traffic into accelerator-friendly execution, improving utilization when latency budgets allow batching. Exact utilization gains depend on model, hardware, arrival rate, and the configured batching window.

⁴ | **NCHW and NHWC (tensor memory layouts):** These acronyms encode the dimension order of 4D image tensors: N (batch), C (channels), H (height), and W (width). In a contiguous NCHW tensor, each channel plane is contiguous; in NHWC, channels for one spatial position are adjacent. Optimized kernels may favor either layout, and a semantic mismatch can scramble values even when the element count is correct. Serving code should therefore make both dimension meaning and any physical-layout conversion explicit.




Figure 13.4: Inference Server Anatomy: A modern inference server decouples network handling from accelerator execution through a staged pipeline. Each stage isolates a concern, from absorbing bursty traffic to forming efficient batches, so the hardware accelerator stays highly utilized despite irregular arrival patterns.

warmup inferences before exposing the model to live traffic. Full registries (versioned artifact stores), release gates (checks before release), and rollback governance (rules for reverting a bad release) belong to chapter 14; the local serving concern is whether the right artifact is loaded and ready. Among these components, the scheduler deserves special attention because it embodies the core serving trade-off between throughput and latency.

The **scheduler** is the “brain” of the inference server. It implements the dynamic batching logic discussed in section 13.7. The scheduler must decide whether to run a single request immediately to minimize its latency or wait five milliseconds for a second request and process them together to maximize throughput.

Systems designers use the **batching window** parameter to tune this trade-off. A window of 0 ms optimizes for pure latency (no batching), while a small bounded window lets the scheduler trade a controlled amount of waiting for higher accelerator utilization. This decision determines how busy the accelerator stays: whether the hardware spends its time computing or waiting for work.

13.3.2 Interface protocols and serialization

The mechanism used to transport data between client and server directly affects the latency budget. Model inference is often highly optimized, yet the cost of moving data into the model (serialization and network protocol overhead) can become the dominant bottleneck, especially for lightweight models where inference time is small.

13.3.2.1 The serialization bottleneck

ML serving payloads are fundamentally different from typical web API payloads: they consist of multi-dimensional float arrays (image tensors, embedding vectors, token ID sequences) that are dense, binary, and large. Text-based formats like JSON are ubiquitous but computationally expensive for this kind of data. Serialization overhead appears when parsing a JSON object requires reading every byte, validating syntax, and converting text representations into machine-native types. For tensor payloads, the cost compounds: floating-point values must first be encoded as variable-length ASCII strings whose size depends on value and formatting, and binary data such as image bytes requires Base64 encoding, which adds 33 percent size overhead before JSON parsing begins. For high-throughput systems, binary alternatives such as Protocol Buffers⁵ and FlatBuffers⁶ preserve CPU cycles for request handling or preprocessing.

⁵ | **Protocol Buffers (Protobuf):** Protobuf uses a predefined schema (from a .proto file) to encode structured data into a compact binary format (Protocol Buffers Authors 2026). Because the schema carries field names and types, the wire payload need not repeat them as JSON does. Its wire format is still not identical to a C++ object’s in-memory layout, so it requires a parsing step and does not provide the same direct zero-copy access pattern that FlatBuffers targets.

⁶ | **FlatBuffers:** The “flat” in the name describes the design: the binary buffer can serve as the serialized representation and the data structure being read, avoiding a separate parsing or unpacking phase for supported access patterns (FlatBuffers Authors 2026). For ML inference, this can enable zero-copy access to tensor metadata—the serving system reads tensor shapes and offsets directly from the buffer rather than allocating a second object representation.

These formats reduce this bloat through schema-aware binary encodings instead of text encodings. Native float arrays can transmit as compact IEEE 754 bytes with no ASCII conversion and no Base64 wrapper. FlatBuffers can also enable zero-copy access in supported cases, where the network buffer can be read without allocating a separate object graph.

13.3.2.2 REST vs. gRPC

Two common interface styles have distinct system characteristics. REST (Representational State Transfer) commonly uses HTTP with JSON. It is widely supported and human-readable, making it a common choice for public-facing APIs. A stateless REST request must carry the information needed to identify or process its state, but it can reference server-side conversation state by an identifier rather than resend an entire large language model (LLM) context. Transport version and payload format are separate choices: REST can use HTTP/2, and an HTTP API need not use JSON. When it does use HTTP/1.1 plus JSON, connection management and text serialization can become visible costs for high-QPS tensor traffic.

In contrast, gRPC (gRPC Remote Procedure Call)⁷ uses HTTP/2 and commonly uses Protobuf. HTTP/2 enables multiplexing multiple requests over a single persistent TCP connection, reducing connection-management overhead and allowing efficient binary streaming. Protobuf provides typed schemas and efficient binary serialization, making gRPC a common choice for internal service-to-service communication where latency and typed interfaces matter.

A concrete payload comparison shows how the serialization choice changes both wire size and parsing cost.

Napkin Math 13.1: JSON vs. Protobuf serialization

Consider a request payload containing 1,000 floating point numbers (for example, an embedding vector).

- **JSON:** Uses ~9 KB on the wire. Requires ~50 μ s to parse.
- **Protobuf:** Uses ~4 KB on the wire. Requires ~5 μ s to parse.

Math: Switching one request to the binary payload saves $50 \mu\text{s} - 5 \mu\text{s} = 45 \mu\text{s}$ of parse time. For this illustrative system processing 10,000 requests per second, the savings compound to $10,000 \times 45 \mu\text{s} = 0.45 \text{ s}$ of CPU time reclaimed every wall-clock second, which is 45 percent of one core freed from serialization overhead alone.

Systems insight: This 10 $\times$ scenario gain makes gRPC/Protobuf, FlatBuffers, or another binary protocol a strong candidate for high-throughput internal microservices when serialization is a visible part of the latency budget.

⁷ | **gRPC (gRPC remote procedure call):** gRPC pairs HTTP/2 transport with an interface definition language and message format, most commonly Protocol Buffers (gRPC Authors 2026). The relevant serving advantage is the combination of typed contracts, persistent multiplexed connections, streaming support, and compact binary messages. The size and latency benefit over REST/JSON depends on payload shape, client/server implementation, and whether serialization is a meaningful share of the end-to-end latency budget.

The system choice is constraint-dependent. REST/HTTP is common when public compatibility, debugging, and ecosystem reach dominate. gRPC/Protobuf, or another binary protocol, is favored when internal high-QPS tensor traffic, connection reuse, or streaming makes serialization a meaningful share of the latency and CPU budget.

The architectural components and protocols examined so far describe *how* serving systems are built. Understanding *why* certain configurations perform better requires analyzing what happens to individual requests as they traverse these components.

13.4 Request Lifecycle

A single HTTP request carrying a 224×224 JPEG image arrives at an inference server. Between the moment the first byte enters the network stack and the moment the classification result leaves, that request traverses six pipeline stages, each consuming milliseconds that the user experiences as wait time. Understanding where time goes within each request is essential for effective optimization: one cannot improve what one does not measure.

13.4.1 The latency budget

For dynamic inference systems, the serving inversion established in section 13.1 creates a latency budget that shapes system design (Gujarati et al. 2020). A serving system with second-scale per-request latency may miss many interactive SLOs, even if it achieves excellent throughput.

Relevant metrics expand from aggregate throughput and mean latency to the latency distribution. Mean and p50 describe the center, while p95 and p99 expose the slow tail. If mean latency is 50 ms but p99 is two seconds, the p99 threshold is 40× the mean, so about 1 percent of requests are at or beyond two seconds. Whether that tail matters depends on the service and its users, which is why percentile targets should come from the SLO rather than convention.⁸

Managing these percentile constraints requires decomposing the total allowed response time into a **latency budget** that allocates time across each processing phase.

⁸ | **Tail latency:** Unlike averages, percentile latencies reveal the performance impact of system outliers common in ML serving, such as model cache misses or garbage collection pauses. These rare, high-latency requests can disproportionately harm user satisfaction and business outcomes. Because the magnitude is workload-specific, services validate percentile targets (p95, p99) against their own users and SLOs.

Definition 13.2: Latency budget

Latency budget is the time allocated to request stages under an end-to-end deadline; the SLO also specifies the required completion quantile.

1. **Significance:** It acts as a zero-sum constraint system where any milliseconds consumed by serialization or network overhead directly reduce the latency budget (L_{lat}) available for model inference.
2. **Distinction:** A latency **SLO** couples a deadline (for example, < 50 ms) with a completion quantile (p99 or p99.9). A service-level indicator (**SLI**) is the measured quantity used to evaluate that objective. Missing an internal SLO consumes reliability margin; crossing an external **SLA** threshold may trigger the remedies defined by the contract. Tail latency spikes can stem from queueing, batch formation, cache state, interference, or data movement.
3. **Common pitfall:** A frequent misconception is that the “model” has the entire budget. In the illustrative breakdown in table 13.3, the remainder is consumed by the request lifecycle (DNS, TLS, load balancing, serialization).

Before computing a full budget, this checkpoint sets the foundational latency-analysis skills every serving engineer needs.

Every serving request decomposes into three phases that each consume part of the latency budget. Preprocessing transforms raw input such as image bytes or text strings into model-ready tensors. Inference executes the model computation. Postprocessing transforms model outputs into user-facing responses.

Checkpoint 13.1: Tail latency and budget

Serving optimizes tail latency under load. Use this checkpoint to separate queueing and batching effects before choosing an optimization.

- ☐ **Queue behavior:** use figure 13.1 to describe why latency rises nonlinearly as utilization approaches saturation.
- ☐ **Tail vs. mean:** can you explain why mean latency cannot demonstrate compliance with a p99 SLO?
- ☐ **Latency budget:** given a 50 ms p99 SLO and 18 ms spent on preprocessing, transfer, and postprocessing, how much remains for inference?

Faster hardware does not automatically mean faster serving. In practice, preprocessing and postprocessing can dominate total latency when inference runs on optimized accelerators. Optimizing exclusively the inference phase yields diminishing returns if the surrounding pipeline remains bottlenecked by CPU operations.

pre infer post

Inference is one slice of the latency budget; preprocessing rivals it.

13.4.2 Latency distribution analysis

Understanding where time goes requires instrumenting each phase independently. A ResNet-50 latency budget breakdown reveals exactly how each millisecond is spent when our classifier receives a JPEG image.

Example 13.4: ResNet-50: Latency budget breakdown

Table 13.3 gives an illustrative ResNet-50 serving request breakdown by phase:

Table 13.3: ResNet-50 latency budget: Illustrative per-phase breakdown of a single serving request, showing that preprocessing and data transfer together rival the cost of the ResNet-50 forward pass itself. The percentages expose where engineering effort pays off in this scenario. Per-phase percentages are rounded to the nearest whole number and may not sum to exactly 100.

Phase	Operation	Time	Percentage
Preprocessing	JPEG decode	3 ms	30%
Preprocessing	Resize to 224×224	1 ms	10%
Preprocessing	Normalize (mean/std)	0.5 ms	5%
Data Transfer	CPU→GPU copy	0.5 ms	5%
Inference	ResNet-50 forward pass	5 ms	50%
Postprocessing	Softmax + top-5	0.1 ms	~1%
Total		10.1 ms	100%

Systems insight: In this illustrative ResNet-50 serving budget, preprocessing consumes 44.6 percent of latency despite model inference being the computationally intensive phase. If TensorRT reduces inference to the assumed 2 ms, preprocessing would dominate at 63.4 percent.

The ResNet example represents compute-bound inference where the forward-pass arithmetic dominates the latency budget. Applying the same framework to a different model architecture often reveals that the bottleneck shifts from compute to memory bandwidth, invalidating the optimization strategies that worked for vision models. Recommendation systems exhibit exactly this shift.

Lighthouse 13.1: Lighthouse example: DLRM serving

Scenario: Serving DLRM with a 10 ms p99 latency budget.

Analysis: While ResNet-50's model stage is dominated by convolutional neural network compute, DLRM's dominant model-stage cost is embedding-table memory access. End-to-end serving bottlenecks still require measuring the full path: preprocessing, inference, postprocessing, and data movement. Table 13.4 breaks the recommendation request down by phase:

Table 13.4: DLRM serving latency: Illustrative per-phase breakdown of a recommendation request under a 10 ms p99 budget, contrasting DLRM's memory-bandwidth-bound embedding lookups against ResNet-50's compute-bound forward pass. Adding compute to the inference stage does not help once embedding-table bandwidth is the binding constraint.

Phase	Operation	Time	Bottleneck
Input Parsing	Request parsing	0.5 ms	CPU
Embedding Lookup	Fetch 100+ dense vectors	6 ms	memory bandwidth
Inference	MLP forward pass	1.5 ms	Compute
Postprocessing	Ranking & Filtering	1 ms	CPU
Total		9 ms	

Systems insight: In DLRM, the “Inference” multilayer perceptron stage is only ~17 percent of the latency. The majority of time is spent in embedding lookups, retrieving massive 128-dim vectors from terabyte-scale tables. This is a memory-bandwidth and capacity-bound workload where adding more compute does not help unless the embedding tables can be served faster.

Together, table 13.3 and table 13.4 expose the same general failure mode: straightforward optimization efforts target where ML expertise applies (model quantization, pruning) while the binding constraint sits elsewhere (image decoding on CPU for ResNet-50, embedding-table memory bandwidth for DLRM). The pattern generalizes: any serving system where the model accounts for less than half of total latency will see diminishing returns from model-only optimizations, regardless of how large those individual speedups are. Amdahl’s Law quantifies the ceiling. Adopting the quantitative approach to serving exposes these hidden bottlenecks before engineering effort is misallocated.

🔍 Systems Perspective 13.3: The quantitative approach to serving

Analysis: Amdahl’s Law at work (section D.2.3 provides the formal derivation): preprocessing (4.5 ms) and data transfer (0.5 ms) consume 49.5 percent of total latency. Optimizing the model $10\times$ faster (5 ms $\rightarrow$ 0.5 ms) yields only $1.8\times$ end-to-end speedup (from 10.1 ms to 5.6 ms). This is why focusing exclusively on model optimization (quantization, pruning) often disappoints: the bottleneck is elsewhere.

Mechanism: DSAs such as TPUs and Tensor Cores replace complex control logic with dense multiply-accumulate arrays, improving achieved throughput for compatible kernels. This makes hardware acceleration an economic requirement for many high-throughput or low-latency serving workloads.

Systems insight: Profile before optimizing. If preprocessing dominates, GPU-accelerated pipelines (NVIDIA DALI) may outperform model quantization.

Moving preprocessing closer to the accelerator can reduce avoidable CPU-GPU transfers, but the end-to-end gain is pipeline-specific. Effective optimization targets the largest time consumers first.

13.4.2.1 The serving tax bill

Beyond the model execution itself, every request pays a “tax” to the serving infrastructure. The relevant overheads span network I/O, serialization, queuing, dispatch, and data movement.

13.4.2.2 The killer microseconds problem

Barroso, Patterson, and colleagues identified a critical gap in how systems handle latency at different time scales (Barroso et al. 2017). Operations in the microsecond range are too short for traditional OS scheduling (which operates at millisecond granularity) yet too long to simply spin-wait without wasting CPU cycles. This “killer microseconds” regime matters in modern serving workloads. Using the representative ranges in table 13.5, serialization at 50–500 μ s, dispatch at 10–50 μ s, and data copy at 100–500 μ s are each individually small, but for a 5 ms inference service, these named microsecond-scale overheads collectively consume about 3.2 percent to 21 percent of the latency budget before network and queuing delays are counted. No single overhead justifies optimization in isolation, yet together they determine whether the system meets its SLO.

Table 13.5: The Serving Tax Bill: A representative breakdown of noninference latency sources. While individual components like serialization can be below 1 ms, they compound. Under the upper ends of the displayed ranges, this tax can rival a 5 ms model invocation. The engineering goal is to measure the actual request path and reduce the binding costs through choices such as binary payloads, persistent connections, and fewer copies.

Tax Component	Typical Cost	Scaling Behavior	Tax Evasion Strategy
Network I/O	1–5 ms	Linear with payload	Compression, Region Colocation
Serialization	50–500 μ s	Linear with payload	gRPC/Protobuf (vs. JSON)
Queuing	0.1–10 ms	Diverges near capacity	Dynamic Batching, Autoscaling
Dispatch	10–50 μ s	Constant per batch	Kernel Fusion (reduce launches)
Data Copy	100–500 μ s	Linear with tensor	Zero-Copy/Shared Memory

The latency budget framework provides a systematic approach to this compound problem. Measurement comes first: without per-phase instrumentation, engineers cannot distinguish a preprocessing bottleneck from a serialization bottleneck, and optimization effort gets misallocated to the

most visible component (the model) rather than the most expensive one. Once measurement reveals the true distribution of time, engineering effort should flow proportionally: a phase consuming 50 percent of latency deserves more attention than one consuming 5 percent, regardless of which feels more tractable. Architectural changes such as GPU-accelerated preprocessing or aggressive batching can shift work between phases entirely, sometimes eliminating a bottleneck rather than merely reducing it.

13.4.3 Resolution and input size trade-offs

Input resolution affects both preprocessing and inference latency, but the relationship differs depending on whether the system is compute bound (limited by arithmetic throughput) or memory-bound (limited by data movement). A compute-bound system slows as computation grows; a memory-bound system slows as transferred bytes grow. Capacity fit alone does not remove bandwidth cost. The roofline analysis in section 11.6 develops this distinction in depth, making it essential for informed resolution decisions.

For convolution-dominated, compute-bound models, equation 13.1 gives a first-order approximation in which throughput scales inversely with resolution squared:

$$\frac{\text{Throughput}(r_2)}{\text{Throughput}(r_1)} = \left(\frac{r_1}{r_2}\right)^2 \quad (13.1)$$

Doubling resolution from 224 to 448 gives a 4× compute-scaling estimate; the illustrative 3.6× scenario includes fixed overhead. Higher resolution can raise convolution arithmetic intensity because weights are reused across more spatial positions, but whether the full model becomes compute-bound depends on the bytes moved by its kernels.

Table 13.6 quantifies how increasing resolution amortizes fixed weight traffic in this simplified model while leaving the executed-kernel bottleneck unresolved.

Table 13.6: Resolution and Modeled Arithmetic Intensity: This estimate divides FLOPs by weight bytes plus one assumed read/write of a scaled activation term. It omits aggregate intermediate-tensor, cache, and workspace traffic, so crossing the V100 ridge point does not establish the executed-kernel bottleneck.

Resolution	Assumed Activation Term	Modeled Arith. Intensity	Bottleneck Classification
224×224	12.5 MB	64.4 FLOP/byte	Not established
384×384	36.7 MB	137 FLOP/byte	Not established
512×512	65.3 MB	183.9 FLOP/byte	Not established
640×640	102.0 MB	218.4 FLOP/byte	Not established

13.4.3.1 Resolution strategies in production

Different deployment contexts impose distinct resolution requirements shaped by their dominant constraints. A mobile application may accept lower resolution (224×224) when latency and battery life outweigh marginal accuracy gains. Some medical-imaging workloads use 512×512 or higher when diagnostic detail warrants the additional compute. A perception pipeline may use lower resolution over a wide field and high-resolution crops for fine-grained recognition. A public API can instead normalize arbitrary uploads to a fixed contract or select among validated resolution profiles. Adaptive resolution is useful only when the selection policy and every chosen profile preserve the task’s quality requirements.

13.4.3.2 Adaptive resolution

Adaptive resolution lets production systems select resolution dynamically based on content. One approach runs a lightweight classifier at 128×128 to categorize content type, then selects task-appropriate resolution with documents at 512×512, landscapes at 224×224, and faces at 384×384. In this illustrative scenario, the policy provides 1.4× throughput improvement with 99.2 percent accuracy retention vs. fixed high resolution. This pattern trades preprocessing cost from running the lightweight classifier for inference savings on the main model.

The latency analysis so far has focused on sequential processing: one request completing before the next begins. The preprocessing, inference, and postprocessing stages use different hardware resources. This separation creates an opportunity to process multiple requests simultaneously.

13.4.4 Hardware utilization and request pipelining

Optimizing each request stage in isolation misses a critical opportunity: the stages use different hardware resources. The latency budget analysis in section 13.4.1 reveals that model inference is only one component of the request lifecycle. From a hardware perspective, the primary goal of a serving system is to maximize the **duty cycle** of the accelerator, the percentage of time the GPU is performing useful computation.

In a serialized serving system, the hardware sits idle during network I/O and CPU-based preprocessing. High-performance serving systems use **request pipelining** to overlap these stages, ensuring the GPU is fed a continuous stream of tensors.

13.4.4.1 Overlapping I/O and compute

The two timing diagrams in figure 13.5 illustrate the impact of pipelining. In the serial case (A), each request must complete its entire lifecycle (Network → CPU Preprocessing → GPU Inference → Postprocessing) before the next request begins, and the grey idle gaps leave the GPU unused for more than 50 percent of the time. In the pipelined case (B), those gaps disappear.




Figure 13.5: Request Pipelining: Pipelining hides latency by overlapping independent operations across different hardware resources. In pipelined execution (B), the CPU processes the next request’s data while the GPU executes the current request’s inference. This can increase accelerator duty cycle and throughput without changing the model.

Pipelining is enabled by asynchronous I/O and concurrency models. Instead of waiting for a GPU kernel to finish, the server’s CPU thread submits the work to the GPU’s command queue and immediately begins preprocessing the next incoming request.

13.4.4.2 The systems metric: Hardware duty cycle

System efficiency measures how fully a serving system saturates its bottleneck resource, usually the GPU’s compute cores or memory bandwidth. Equation 13.2 defines this hardware duty cycle.

$$\text{System Efficiency} = \frac{\sum T_{\text{compute}}}{\text{Wall Clock Time} \times \text{Resource Count}} \quad (13.2)$$

If a ResNet-50 request takes 10 ms total (5 ms GPU, 5 ms CPU), a serial system achieves only 50 percent efficiency. By pipelining just two requests, efficiency approaches 100 percent (assuming the CPU can keep up with the GPU). If the CPU is too slow to feed the GPU, the system becomes CPU-bound, and further model optimization provides zero throughput gain. This is Amdahl’s Law from section D.2.3 applied to serving: if preprocessing consumes 50 percent of latency, maximum speedup is $2\times$ regardless of how fast the model runs. The hardware trajectory makes this ceiling progressively tighter. Accelerator compute throughput (FLOPs) has grown far faster than CPU single-thread performance across successive hardware generations, so the inference portion of the pipeline shrinks while the CPU-bound preprocessing portion remains unchanged. A system that was compute-bound on an older accelerator may become CPU-bound after a hardware upgrade—not because preprocessing got slower, but because the model got dramatically faster while the CPU did not.

13.4.5 Postprocessing

The request lifecycle concludes with postprocessing, the phase that transforms model outputs into actionable results. A neural network produces raw tensors (floating-point arrays that carry no

inherent meaning to applications or users). A 0.95 probability becomes a confident “dog” label only after postprocessing converts it; a sequence of token IDs becomes readable text; a bounding box tensor becomes a highlighted region in an image. Postprocessing significantly impacts both latency and the usefulness of predictions.

13.4.5.1 From logits to predictions

Classification models output logits or probabilities across classes. Converting these raw outputs to predictions involves several steps. The simplest is argmax selection, which returns the highest-probability class. Thresholding applies a confidence cutoff, returning predictions only when the model is sufficiently certain. Top- k extraction returns multiple high-probability classes with their scores, useful when applications need ranked alternatives. Calibration adjusts raw probabilities to better reflect true likelihoods, a step that adds computation but is essential when downstream systems make decisions based on confidence scores. For ResNet-50 image classification, listing 13.1 shows the full postprocessing path from raw logits to an API-ready response, including probability normalization, top- k extraction, label lookup, and response formatting.

Listing 13.1: ResNet-50 Postprocessing: Normalizes raw logits into probabilities, extracts top- k predictions, and formats the API response.

```
# Normalize raw logits into probabilities
# Input: logits tensor of shape (1, 1000), one score per ImageNet
# class
probs = torch.softmax(logits, dim=-1) # Normalize to sum=1 on GPU

# Extract top-5 predictions for multi-class response
# topk returns (values, indices) sorted by probability
top5_probs, top5_indices = probs.topk(5) # GPU operation
top5_probs = top5_probs.squeeze(0).tolist()
top5_indices = top5_indices.squeeze(0).tolist()

# Map class indices to human-readable labels
# imagenet_labels: list of 1000 class names from synset mapping
labels = [imagenet_labels[i] for i in top5_indices] # CPU lookup

# Format response with predictions and metadata for API contract
response = {
    "predictions": [
        {"label": label, "confidence": float(prob)}
        for label, prob in zip(labels, top5_probs)
    ],
    "model_version": "resnet50-v2.1", # Client-side version tracking
    "inference_time_ms": 5.2, # Observability for latency monitoring
}
```

For this example, total postprocessing time is approximately 0.1 ms, negligible compared with preprocessing and inference. Each step adds latency but improves response utility. Fitting a calibration mapping is an offline task; applying a simple fitted mapping online is usually inexpensive, but it is valuable only when downstream decisions require confidence scores whose empirical meaning has been validated.

13.4.5.2 Output formatting

Production systems rarely return raw predictions. Outputs must conform to API contracts that specify JSON serialization schemas, confidence score formatting, and thresholding rules. Error handling must address edge cases: the system must define behavior when no prediction exceeds the confidence threshold or when the input appears out-of-distribution. Response metadata (model version, inference time, feature attributions) enables downstream monitoring and debugging.

The latency budget analysis reveals where time goes within a single request. Production systems, however, do not process requests in isolation: they must handle hundreds or thousands of concurrent requests competing for finite resources. Understanding this concurrency requires a different analytical framework.

Tracing one request identifies the work and latency in each stage; once requests overlap, those stages compete for finite resources and queuing delay enters the budget.

13.5 Queuing Theory

In production, concurrent requests compete for finite resources, and queuing theory predicts how this competition affects latency. These principles explain the counterintuitive behavior that causes well-provisioned systems to violate latency SLOs when load increases modestly.

13.5.1 Little's Law

Serving engineers routinely face a concrete capacity decision: given an expected request rate and measured mean response time, the system must determine how much work is in flight before deciding how many GPUs to provision. Little's Law (section D.2.4) answers the first question by relating mean population to throughput; a tail-latency SLO can only provide a planning proxy. The M/M/1 model later answers the second by predicting how latency degrades under load. Together, they provide the quantitative framework for capacity planning.

🔍 Systems Perspective 13.4: Notation alert: L vs. latency

In queuing theory, T_{lat} denotes the *response time* or *time in system per request*; queue-only waiting time is W_q . Here, Q_{req} denotes the average in-system request count, λ_{arr} for arrival rate, and $\rho_{\text{serv}} = \lambda_{\text{arr}}/\mu$ for serving utilization. The subscripts distinguish queueing notation from the degradation equation's λ sensitivity parameter and keep serving utilization from occupying bare ρ . In subsequent latency-budget equations (such as equation 13.4 and equation 13.7), descriptive $L_{\text{lat},*}$ terms name components of the budget, such as waiting and compute. In the batching analysis that follows (section 13.7.3), $L_{\text{lat},\text{wait}}$ corresponds to the queueing wait component W_q , and $L_{\text{lat},\text{compute}}$ includes inference time.

Serving engineers need a tool that connects observable metrics to capacity requirements. The most celebrated result in queuing theory is **Little's law**,⁹ which equation 13.3 expresses as a simple relationship between three quantities in any stable system:

$$Q_{\text{req}} = \lambda_{\text{arr}} \cdot T_{\text{lat}} \quad (13.3)$$

where Q_{req} is the average number of requests in the system, λ_{arr} is the arrival rate (requests per second), and T_{lat} is the average time each request spends in the system.

Concretely, for a target arrival rate of 1000 QPS and a mean response time of 50 ms, Little's Law translates that pair into the average in-system population; it does not set a minimum GPU batch size or a hard activation-memory floor. The following worked example carries out that calculation.

This relationship holds regardless of arrival distribution, service time distribution, or scheduling policy. A practical capacity calculation shows why this universality matters for serving memory.

📝 Napkin Math 13.2: Little's Law capacity sizing

Problem: How many requests are in the system on average at 1,000 QPS?

Math: Little's Law gives $Q_{\text{req}} = \lambda_{\text{arr}} T_{\text{lat}}$, so average population equals throughput multiplied by mean time in system (section D.2.4 derives the law).

Given:

- **Throughput target** (λ_{arr}): 1,000 QPS.
- **Mean time in system** (T_{lat}): 50 ms (0.05 s).

Math: $Q_{\text{req}} = 1,000 \text{ QPS} \times 0.05 \text{ s} = 50$ concurrent requests on average

Systems insight: The system holds 50 requests on average across batch and queue state. A GPU batch-size limit of 32 does not by itself make 1,000 QPS impossible because queued requests

⁹ | **Little's law:** John D. C. Little (1961) proved that $Q_{\text{req}} = \lambda_{\text{arr}} T_{\text{lat}}$ holds for stable long-run averages over consistent system boundaries, regardless of arrival distribution, service distribution, or scheduling discipline; this universality anchors ML capacity planning: the formula requires no Poisson-arrival or exponential-service assumption. It does require well-defined long-run averages and consistent definitions of arrival rate, time in system, and population.

may reside in host memory; service rate, queueing behavior, and the tail distribution determine feasibility.

Little's Law has immediate practical implications. If an inference service averages 10 ms per request ($T_{\text{lat}} = 0.01$ s) and the system shows 50 concurrent requests on average ($Q_{\text{req}} = 50$), then the arrival rate must be $\lambda_{\text{arr}} = Q_{\text{req}}/T_{\text{lat}} = 5000$ requests per second. Conversely, if the system limits the average in-system population to 10 and measures 10 ms average time in system, the corresponding throughput is 1,000 requests per second.

13.5.2 The batching tax: The latency-throughput frontier

While Little's Law relates queue depth to throughput, it does not account for the **batching tax**: the deliberate delay introduced to maximize hardware utilization. This delay creates a queueing problem.

When an inference server uses request-triggered fixed-size batching, it introduces two distinct sources of latency. **Batch formation delay** ($L_{\text{lat,form}}$) is the time each request waits for the batch to fill; the first request waits longest, and the last waits zero. **Inference inflation** is the growth in inference time $T_{\text{inf}}(B)$ when the GPU processes B samples instead of 1. The resulting latency-throughput Pareto frontier is the set of configurations where one cannot improve throughput without paying a “tax” in increased latency. Under stationary Poisson arrivals, equation 13.4 approximates the average per-request latency for batch size B and arrival rate λ_{arr} :

$$L_{\text{lat,total}} \approx \underbrace{\frac{B-1}{2\lambda_{\text{arr}}}}_{\text{Formation delay}} + \underbrace{T_{\text{inf}}(B)}_{\text{Inference time}} \quad (13.4)$$

Equation 13.4 reveals the “cost of throughput.” Increasing B to saturate the GPU amortizes the hardware cost, but inflates the per-request latency. Concretely, at 500 QPS, moving from batch-1 to batch-32 increases wait-time from 0 ms to 31 ms, contributing to a 23× total latency penalty (2 ms → 46 ms). For a systems engineer, this tax is the primary regulator of economic efficiency: the engineer chooses the batch size that maximizes throughput (minimizing cost per query) without violating the latency SLO (L_{lat}).

13.5.3 The utilization-latency relationship

Little's Law describes average system behavior, but it does not reveal *how* latency changes as load approaches capacity. The M/M/1 queue model answers the critical question of how much spare capacity a serving system needs (Harchol-Balter 2013).¹⁰ For a system with Poisson arrivals and exponential service times, equation 13.5 gives the average time in system:

$$T_{\text{lat}} = \frac{1}{\mu - \lambda_{\text{arr}}} = \frac{\text{service time}}{1 - \rho_{\text{serv}}} \quad (13.5)$$

where λ_{arr} is the arrival rate, μ is the service rate (requests per second the server can handle), and $\rho_{\text{serv}} = \lambda_{\text{arr}}/\mu$ is the utilization (fraction of time the server is busy).

Equation 13.5 reveals why serving systems exhibit nonlinear behavior: small increases in load near capacity cause disproportionate latency increases.¹¹ Table 13.7 quantifies this relationship, showing how average time in system grows rapidly as utilization approaches 100 percent.

The M/M/1 model assumes exponentially distributed service times, whereas fixed-shape inference at a fixed batch size can have lower service-time variance. An M/D/1 model can then approximate the mean queueing wait more closely. This chapter uses M/M/1 because it yields closed-form solutions and preserves a conservative margin when its assumptions are stated. Real services still require measured arrival and service-time distributions, especially when input shapes, batching, preprocessing, or cache state vary.¹²

10 | M/M/1 queue: Queueing theory originated with Agner Krarup Erlang's 1909 analysis of the Copenhagen Telephone Exchange, where call arrivals were modeled as memoryless (Poisson). The M/M/1 model's exponential service-time assumption introduces more variance than many fixed-shape ML inference workloads exhibit. At the same utilization and mean service time, its mean *queueing wait* is twice the M/D/1 value; the ratio for total response time approaches two only as utilization becomes high. The model is therefore a useful conservative approximation, not a universal description of inference traffic.

11 | Super-linear latency divergence: The planning knee often appears well before full saturation; in the M/M/1 mean response-time equation, $E[T] = \frac{1/\mu}{1 - \rho_{\text{serv}}}$, where $\rho_{\text{serv}} = \lambda_{\text{arr}}/\mu$ is utilization. The $(1 - \rho_{\text{serv}})^{-1}$ term diverges as $\rho_{\text{serv}} \rightarrow 1$: at $\rho_{\text{serv}} = 0.7$, mean response time is already 3.3× the base service time; at $\rho_{\text{serv}} = 0.9$, it is 10×. The exact operating limit is a policy and workload choice, but trying to stretch a latency-sensitive queue toward saturation creates disproportionate latency growth.

12 | Kendall notation: In the A/S/c (arrival/service/servers) notation, “M” signifies a Markovian (memoryless) process and “D” means deterministic. M/M/1 is conservative relative to M/D/1 for mean queueing wait when service times are nearly deterministic, but neither model replaces measurement of a production workload's arrival process, service-time variance, or tail behavior.

Table 13.7: Utilization-Latency Relationship: Average time in system (wait + service) as a multiple of service time for an M/M/1 queue. At 50 percent utilization, time in system is 2× service time; at 90 percent, it reaches 10×. This nonlinear growth explains why systems that perform well at moderate load suddenly violate SLOs when traffic increases: moving from 80 percent to 90 percent utilization doubles latency.

Utilization (ρ_{serv})	Latency Multiple	Example (5 ms service)
50%	2×	10 ms
70%	3.3×	17 ms
80%	5×	25 ms
90%	10×	50 ms
95%	20×	100 ms

13.5.4 Multi-server considerations

The single-node queuing analysis in this section focuses on a single serving node (one machine serving inference requests). This scope focuses on mastering the basic unit of ML systems. Single-node queuing dynamics are prerequisite to effective scaling. Engineers cannot optimize a distributed system without first understanding the behavior of its components.

M/M/1 analysis provides a first approximation for right-sizing individual nodes, identifying a scaling trigger, and avoiding premature scale-out. It can test whether one modeled server has enough headroom at an expected arrival rate and show how latency grows as modeled load approaches capacity. Measurements must then determine whether the actual bottleneck is node capacity, batching policy, preprocessing, cold start, runtime configuration, or a departure from the queue model’s assumptions.

Once traffic truly exceeds single-node capacity, the next move is replica-level scale-out: multiple independent serving nodes sit behind a load balancer and each runs the same model. The M/M/c queuing model extends M/M/1 to c parallel servers, showing how replicas can improve latency when traffic is balanced across independent servers. The exact p99 improvement depends on arrival process, service-time variance, dispatch policy, and per-replica utilization. That replica model is still different from distributed inference, where one request is split across GPUs through model sharding, tensor parallelism, or pipeline parallelism. This chapter establishes the single-node and replica foundations; distributed inference adds coordination overhead and consistency challenges beyond this scope.

13.5.5 Tail latency

Production SLOs typically specify percentile targets (p95, p99) rather than averages because tail latency determines user experience for the slowest requests (Dean and Barroso 2013). For an M/M/1 queue, the p99 latency follows:

$$T_{\text{lat,p99}} \approx \frac{\text{service time}}{1 - \rho_{\text{serv}}} \cdot \ln \left(\frac{1}{1 - 0.99} \right) \approx \frac{4.6 \cdot \text{service time}}{1 - \rho_{\text{serv}}} \quad (13.6)$$

At 70 percent utilization, the M/M/1 p99 approximation is approximately 15 times the service time ($4.6/0.3 \approx 15.3$), while average latency is only 3.3 times. For deterministic-service models such as M/D/1, tail values require model-specific calculation rather than a simple universal multiplier. The important point is unchanged: systems that seem healthy with low average latency can have unacceptable tail latency, since the average hides the experience of the unluckiest requests.

13.5.5.1 The tail at scale problem

Dean and Barroso’s analysis reveals *why* tail latency becomes critical as systems scale beyond single machines (Dean and Barroso 2013). When requests fan out to multiple servers, the probability of experiencing at least one slow response grows rapidly with server count. This “tail at scale” effect makes individual server tail latency critical for overall system performance.

For single-machine serving, this principle has two implications. First, tail latency on individual machines matters because it will compound when systems eventually scale. Second, the tail-tolerant

techniques described in section 13.5.6 (hedging, graceful degradation) provide value even on single machines and become indispensable at scale.

Tail-tolerant techniques such as request hedging send redundant requests after a timeout, accepting whichever response arrives first. Backup requests and load balancing away from slow servers directly address latency variance. These techniques apply cleanly to multiple model replicas, and some single-node systems can approximate them with concurrent streams or instances when cancellation and resource isolation semantics permit it. They become essential when scaling to distributed inference systems.

The queuing model and tail latency analysis provide the inputs for capacity planning. A concrete deployment makes the trade-offs tangible.

Applying the M/M/1 tail-response-time model to ResNet-50 makes the capacity constraint concrete.



Napkin Math 13.3: ResNet-50 capacity planning

Consider designing a ResNet-50 serving system with these requirements:

- **Target p99 latency:** 50 ms
- **Peak expected traffic:** 5,000 QPS
- **Service time** (TensorRT FP16): 5 ms

Step 1: Find safe utilization. From equation 13.6, $T_{lat,p99} \approx 4.6 \times \text{service time} / (1 - \rho_{serv})$. Setting $T_{lat,p99} \leq 50$ ms with 5 ms service time gives $\rho_{serv} \leq 1 - (4.6 \times 5 \text{ ms}) / 50 \text{ ms} = 0.54$ (54 percent maximum utilization). This uses the conservative M/M/1 p99 bound from the displayed equation rather than applying an average-wait M/D/1 adjustment to a tail-latency SLO.

Step 2: Calculate required service rate. $\mu_{required} = 5,000 \text{ QPS} / 0.54 = 9259.3 \text{ req/s}$

Step 3: Determine GPU count. In this M/M/1 model, the single-GPU rate is 200 req/s.

GPUs needed = $9259.3 \text{ req/s} / 200 \text{ req/s} = 46.3 \rightarrow 47$ GPUs

Step 4: Add headroom for variance. This scenario adds 30 percent headroom for traffic spikes and variance: final count = $47 \times 1.3 = 61.1$, rounded up to 62.

Step 5: Verify fault tolerance. The 30 percent headroom addresses traffic variance, but production systems also need fault tolerance. With 62 GPUs, losing one leaves 61 GPUs handling 5,000 QPS. The postfailure utilization is $5,000 \text{ QPS} / (200 \text{ req/s} \times 61) = 41.0\%$.

This remains well below the 54 percent safe utilization threshold, confirming N+1 redundancy is satisfied. For stricter fault tolerance requirements, N+2 redundancy (tolerating two simultaneous failures) would require 49 GPUs under the same safe-utilization threshold, or about 64 GPUs if the 30 percent headroom must remain after two simultaneous failures.

Result: Provision 62 V100 GPUs to serve 5,000 QPS at 50 ms p99 latency with N+1 fault tolerance.

The queuing analysis explains the capacity planning approach detailed in section 13.11.3 and connects directly to the MLPerf Server scenario. Section 12.8.4.2 explains that an MLPerf Server result is valid only when the run satisfies its percentile-latency constraint; otherwise, the target QPS must be reduced and the run repeated.

13.5.6 Tail-tolerant techniques

Eliminating all sources of latency variability is often impractical. Production systems instead employ techniques that tolerate variability while still meeting SLOs (Dean and Barroso 2013; Dean 2012). The useful organization is by failure mode: a straggler replica calls for a race, fan-out calls for early detection, overload calls for admission control or graceful degradation, and retry amplification calls for coordinated shedding.

For a straggler replica, the system can race the slow path. Under hedging, when a request has not completed within the expected time, the system sends a duplicate request to another server.¹³ The client uses whichever response arrives first and cancels the other. For ML serving, this means

¹³ | **Hedging:** The term is borrowed from finance, where an offsetting bet reduces risk; here, the redundant request is a bet against a slow server. This is not free: for ML systems, the losing hedged request may still occupy accelerator time if inference has already launched, because ordinary GPU kernels are not cheaply cancelled mid-execution. Thus, a hedging policy must budget duplicate work as well as the latency benefit.

maintaining multiple model replicas and routing slow requests to alternative replicas. With a stable latency distribution, a threshold set at the historical 95th percentile would trigger duplicates for about 5 percent of requests; correlated slowdowns can raise that fraction and make hedging amplify overload instead of reducing the tail.

Ordinary launched inference kernels are not cheaply interrupted mid-execution. When a hedged request completes, the duplicate must be cancelled, but if inference has already begun on the GPU, cancellation approaches include checking a cancellation flag before launching inference, accepting wasted compute for the in-flight kernel, or using request prioritization to deprioritize the duplicate. Since hedging typically applies only to a small tail of requests, the overhead from occasional wasted compute can remain acceptable when the policy is tuned carefully.

Tied requests make the same race more aggressive by placing the request in multiple server queues simultaneously, with coordination that cancels the other copies once one server begins processing. This avoids waiting to detect a slow response before hedging and can reduce queueing delay when multiple ready replicas are available. It also creates duplicate queue entries, so the coordination mechanism must cancel losing copies before they consume scarce accelerator time.

Fan-out systems need a different intervention point because one slow backend can stall the entire distributed request. Canary requests first send the request to a small subset of one to two servers.¹⁴ If these return within expected time, the system sends to the remainder. If the canary is slow, the system can retry elsewhere or use cached results before committing to the full fan-out. The technique turns a potential tail-latency amplification problem into an early warning signal.

When the problem is overload rather than a single straggler, racing makes the system worse by adding duplicate work. The system instead has to protect user-visible responsiveness and admitted-request latency. Graceful degradation returns approximate results rather than timing out: classification systems can return cached predictions for similar inputs, generative models can return shorter outputs, and ensembles can return predictions from a subset of models. Reducing the number of active ensemble members during overload directly shortens service time (T_{svc}), which increases the server's service rate μ and brings utilization $\rho_{\text{serv}} = \lambda_{\text{arr}}/\mu$ back below the queueing knee (figure 13.1), trading a controlled accuracy reduction for SLO survival instead of an uncontrolled latency collapse. Admission control is stricter. When queue depth exceeds a threshold, it proactively rejects requests with immediate 503 responses rather than accepting work that is likely to time out. This sacrifices throughput to protect latency for admitted requests.

¹⁴ | **Canary:** Named for the coal mine practice (early 1900s–1980s) of using birds whose high metabolic rate made them sensitive to toxic gases before concentrations became lethal to humans. In ML serving, canary requests serve the same early-warning function for fan-out queries: by testing 1–2 backends before committing to the full fan-out, the system detects slow or failing replicas before a single straggler stalls the entire distributed inference request—a critical protection when fan-out width means tail latency grows with the maximum of all backend response times.



Checkpoint 13.2: Queuing and SLO headroom

Latency SLOs are not enforced by *fast inference* alone; they are enforced by *headroom*.

- ☐ **Little's Law:** can you use $Q_{\text{req}} = \lambda_{\text{arr}} T_{\text{lat}}$ to explain why rising queue depth implies rising latency even if per-request compute time is unchanged?
- ☐ **Utilization cliff:** can you explain why latency grows nonlinearly as utilization ρ_{serv} approaches one, and why production systems target a conservative ρ_{serv} rather than “100 percent busy”?
- ☐ **Wait vs. compute:** given an end-to-end latency budget, can you separate $L_{\text{lat,compute}}$ from $L_{\text{lat,wait}}$ and explain which one queuing theory primarily predicts?
- ☐ **Capacity planning:** can you explain why a throughput number is only “real” if requests still meet the percentile latency SLO under load?
- ☐ **Headroom estimate:** to hold p99 under 50 ms at 2,000 requests per second, estimate the average in-flight request count $Q_{\text{req}} = \lambda_{\text{arr}} T_{\text{lat}}$ implies, and how fast that margin shrinks as ρ_{serv} passes 0.7.

One illustrative starting point is a queue containing two to three service quanta per worker; with four workers, that corresponds to 8 to 12 queued requests. This is a hypothesis to test, not a portable threshold. Adaptive admission control can tighten the limit when observed p99 latency rises above target and relax it when measured headroom remains. Fixed-shape compiled inference sometimes has a narrower service-time distribution than a general web request, which can make the threshold

easier to estimate. Variable inputs, batching, cache state, preprocessing, and shared resources can remove that advantage, so production controllers still need measured distributions rather than an assumed $M/D/1$ queue.

A subtle failure mode occurs when all replicas are overloaded simultaneously. If the load balancer retries rejected requests at other replicas that are also overloaded, retry traffic amplifies the overload. Coordinated load shedding addresses this by sharing load information across replicas, enabling system-wide decisions about which requests to accept. When global load exceeds capacity, replicas collectively reject the same fraction of requests rather than each rejecting independently and triggering retries.

These techniques become essential at scale when fan-out amplification makes individual server tail latency visible to users. Single-machine serving systems can implement hedged and tied requests across GPU streams or model replicas. The queuing analysis here assumes first-in-first-out (FIFO) processing, but production systems often implement priority scheduling such as deadline-aware or shortest-job-first approaches to further reduce tail latency for heterogeneous workloads (Harchol-Balter 2013).

The tail-tolerant techniques examined here optimize the flow of requests through a functioning serving system. The queuing analysis, however, assumes two critical preconditions: that models are loaded and ready to process requests, and that predictions match what was validated during development. In production, this assumption fails regularly: during deployments, new instances must load models from scratch; during scaling events, cold start latency affects the first requests to new replicas; and when preprocessing pipelines diverge from training, accuracy silently degrades. Section 13.6 examines these lifecycle challenges that must be solved before queuing optimization becomes relevant.

13.6 Model Lifecycle Management

Queuing theory and tail-tolerant techniques optimize the steady-state flow of requests, but they cannot help if the system never reaches steady state. In an illustrative deployment, a newly deployed replica takes 35 seconds to compile its TensorRT engine, while a PIL/OpenCV preprocessing mismatch costs five percentage points of accuracy. Addressing such lifecycle failures requires engineering discipline in two areas: getting models ready to serve (cold start and initialization) and keeping predictions faithful to what was validated (training-serving skew).

13.6.1 Training-serving skew

A model that performed well during validation may silently degrade when deployed. This phenomenon, known as training-serving skew, represents one of the most subtle failure modes in production ML because it is invisible to latency monitoring and exception tracking (Sculley et al. 2015; Baylor et al. 2017).

Definition 13.3: Training-serving skew

Training-serving skew is the distributional divergence between the training and inference environments caused by inconsistent logic or state.

1. **Significance:** It violates the consistency imperative and can cause silent accuracy degradation when the transformation functions differ ($f_{\text{train}}(x) \neq f_{\text{serve}}(x)$).
2. **Distinction:** Unlike data drift (which is an external shift in the environment), training-serving skew is an internal failure of the engineering stack.
3. **Common pitfall:** A frequent misconception is that skew is “found” by looking for errors. In reality, it is invisible to exceptions: the system runs perfectly and the latency is low, but the predictions are statistically wrong.

Section 14.4.1.2 develops skew diagnosis, monitoring, and organizational prevention. Its serving-specific manifestation is **preprocessing divergence**: the real-time inference pipeline processes raw data differently than the batch training pipeline, a common failure mode when training uses

Python/Pandas while serving uses C++/Java or optimized inference servers. Unlike data drift, this divergence is internally introduced and often preventable through shared transformations, parity tests, and careful engineering.



Example 13.5: ResNet-50: Image preprocessing skew

Scenario: ResNet-50 vision serving models experience silent prediction accuracy degradation between offline training and online serving pipelines.

Diagnosis: The serving pipeline implemented C++/OpenCV preprocessing (cv2.INTER_LINEAR, BGR order) while training used Python/PIL (PIL.BILINEAR, RGB order, different normalization constants).

Systems lesson: Preprocessing divergence causes silent training-serving skew without raising runtime exceptions. Standardizing data transformation pipelines across training and serving environments prevents input distribution corruption.

13.6.2 Cold start and initialization dynamics

With preprocessing pipelines designed to avoid training-serving skew, the next challenge is getting models ready to serve. Before processing any request, models must load from storage into memory and prepare for inference (Romero et al. 2021). This initialization latency, known as cold start, affects system responsiveness during deployments, scaling events, and recovery from failures.



Definition 13.4: Cold start

Cold start is the initialization latency incurred when instantiating a new model replica.

1. **Significance:** It represents the fixed cost of state hydration (loading weights, compiling graphs), which can take seconds or minutes, effectively blocking the system's ability to scale elastically in response to traffic bursts.
2. **Distinction:** Unlike inference latency (L_{lat}), which is a per-request cost, cold start is an initialization cost paid when a replica or model instance is created, including deployment, scale-out, recovery, and on-demand model loading.
3. **Common pitfall:** A frequent misconception is that cold start is “just loading weights.” In reality, it includes graph compilation and memory allocation, which can often take longer than the bandwidth-limited data transfer itself (D_{vol}/BW).

Cold start dynamics determine whether systems meet latency requirements from the moment they begin serving traffic. Breaking a representative startup into phases reveals where each part contributes to total initialization latency.

Cold start latency compounds from multiple sources, each adding to the time between deployment and serving readiness. **Weight loading** reads model parameters from disk or network storage. Graph compilation performs just-in-time compilation of operations for the specific hardware. Memory allocation reserves GPU memory for activations and intermediate values. Warmup¹⁵ execution performs initial inferences that populate caches and trigger lazy initialization.

The CUDA context¹⁶ is the first cost in the cold start timeline. Before any GPU operation, the CUDA runtime must establish a *context*: a data structure that tracks memory allocations, loaded kernels, and device state. Creating a context requires communicating with the GPU driver and allocating GPU memory for internal bookkeeping. The 0.3–0.5 s value in table 13.8 is a scenario assumption for this cold-start budget, not a universal CUDA constant. CUDA lazy loading defers some module and kernel loading until first use, reducing apparent startup time but shifting some cost to the first inference (NVIDIA 2026a).

CUDA MPS (Multi-Process Service)¹⁷ addresses GPU sharing for multi-process deployments. Normally, each process creates its own CUDA context, and the GPU may time-slice between contexts. MPS allows work from multiple processes to overlap on the GPU through a shared service, reducing context-switching overhead and improving utilization when individual processes underuse the

¹⁵ | **Warmup:** Borrowed from just-in-time (JIT) compilation, where initial executions compile hot paths into optimized machine code; for ML serving, warmup inferences trigger CUDA kernel compilation, cuDNN algorithm auto-tuning, and memory pool allocation that frameworks defer until first use. Without warmup, the first live request absorbs all of this setup and can be orders of magnitude slower than steady state. During autoscaling events, this means new replicas can violate SLOs for their first several seconds of traffic.

¹⁶ | **CUDA (compute unified device architecture):** NVIDIA's parallel computing platform (Nickolls et al. 2008), named for its goal of unifying diverse GPU shader models into a single general-purpose architecture. Before CUDA, GPU programming required disguising computations as graphics operations. The CUDA context—the data structure tracking memory allocations, loaded kernels, and device state—is the runtime's per-process gateway to GPU resources; in serverless or rapidly scaling serving systems, context creation and lazy module loading can become visible parts of cold start latency (NVIDIA 2026a).

¹⁷ | **CUDA MPS (multi-process service):** MPS creates a control daemon that allows CUDA work from different processes to overlap on the GPU, which can improve utilization and reduce context-switching overhead when processes individually underuse the accelerator (NVIDIA 2026d). For multi-model serving, MPS can help replicas share GPU streaming multiprocessors efficiently. The trade-off is fault isolation: clients share MPS-managed GPU state, so hardware partitioning with Multi-Instance GPU (MIG) provides stronger isolation at the cost of fixed partition granularity.

device (NVIDIA 2026d). The trade-off is reduced isolation: a crash in one process can affect others sharing the MPS server.

Example 13.6: ResNet-50: Cold start timeline

Table 13.8 traces the per-phase duration of a ResNet-50 cold start:

Table 13.8: ResNet-50 cold start timeline: Illustrative per-phase durations for weight loading, CUDA context creation, TensorRT compilation, and warmup, with totals for the optimized local case and the first-deploy cloud case showing where the dominant cost lives.

Phase	Duration	Notes
Weight loading (SSD)	0.5 s	98 MB FP32 weights from local storage
Weight loading (S3)	3–5 s	Network latency dominates for cloud storage
CUDA context	0.3–0.5 s	GPU driver initialization and memory setup
TensorRT compilation	15–30 s	Converts PyTorch model to optimized engine
Warmup (10 inferences)	0.2 s	Triggers remaining lazy initialization
Runtime overhead	0.4 s	Process startup, framework hooks, and runtime setup
Total (local, optimized)	~1.5 s	With precompiled TensorRT engine, warm container
Total (cloud, first deploy)	~35 s	Including compilation from cold state

Systems insight: Precompiling models and storing the optimized engine eliminates the 30-second compilation phase on subsequent deployments.

Without warmup, the first real request triggers compilation and memory allocation mid-inference, often causing timeout failures. A request that normally takes 5 ms might require 500 ms during cold start, violating SLOs and degrading user experience.

13.6.3 Loading strategies

Different loading strategies trade off cold start duration against serving performance and memory efficiency. The simplest approach, full loading, reads the entire model into memory before serving begins. This maximizes inference speed since all weights are immediately available, but extends cold start duration and limits model size to available memory. The approach is appropriate when cold start latency is acceptable and models comfortably fit in memory.

Memory mapping offers an alternative by mapping model files into the process address space and faulting pages in on access. It can reduce eager copying and shorten apparent startup, but a dense model will eventually touch nearly all of its weights, so deferred page faults can move loading latency into early requests. Selective benefit is greater for sparsely accessed components or model repositories that load subsets on demand; latency-sensitive services still prefault or warm the pages needed on the request path.

A third strategy, lazy initialization, defers compilation and allocation until first use. This minimizes startup time but shifts latency to the first request. Production systems often combine lazy initialization with synthetic warmup requests to trigger initialization before real traffic arrives.

13.6.4 Model caching infrastructure

Production systems cache model weights at the infrastructure level to reduce cold start for common deployment scenarios. One approach, container image embedding, bundles model weights directly in the container image. This produces a single deployment artifact and eliminates network fetches at startup, but creates large images (often 10–50 GB) that slow container pulls and consume registry storage. This approach works best for models that rarely update.

For organizations with many models and frequent updates, a shared filesystem (EFS, GCS FUSE) containing versioned model weights provides a more flexible alternative. Multiple replicas can read the same artifacts without baking each version into a container. Updates still require an explicit, atomic version switch and rollout; allowing files to change underneath live replicas would undermine

reproducibility. The trade-off is that network latency affects cold start and filesystem availability becomes a critical dependency.

When cold start latency is critical for high-traffic models, node-local SSD caching prepopulates local SSDs on inference nodes with frequently-used models. This approach provides fast loading from Non-Volatile Memory Express (NVMe) drives at 500 MB/s or more without network dependency, but requires cache management to handle model updates and capacity limits. The choice among these strategies depends on model update frequency: infrequent updates favor container embedding, frequent updates favor shared filesystem, and performance-critical deployments benefit from local caching with background refresh.

13.6.5 Multi-model serving

Production systems often serve multiple models from a single machine, whether different model versions for A/B testing, ensemble components, or entirely different models sharing infrastructure. GPU memory becomes the limiting resource, requiring careful management strategies.

Three strategies address multi-model memory management. Time-multiplexing loads one model at a time and swaps based on request routing; this approach is simple but introduces swap latency. Memory sharing partitions GPU memory among models, limiting concurrent execution count but enabling more models to remain resident. Model virtualization, as implemented by frameworks like Triton, separates model lifecycle from application code through model repository and control APIs for loading, unloading, and versioning models (NVIDIA 2024d, 2026c, 2026b). The choice depends on request patterns: if models receive traffic evenly, concurrent loading works; if traffic is bursty and model-specific, time-multiplexing with explicit preloading reduces average latency while maximizing GPU utilization.

13.6.5.1 Multi-stream execution

When multiple models or multiple instances of the same model must run concurrently on a single GPU, the hardware must partition resources between them. NVIDIA's Multi-Instance GPU¹⁸ technology enables hardware-level isolation, dividing an A100 into up to seven independent GPU instances, each with dedicated memory and compute resources (NVIDIA 2026e). MIG is available on A100, A30 (up to four instances), H100, H200, and newer data center GPUs. For older GPUs such as V100 or T4, CUDA stream scheduling provides time-multiplexed sharing without hardware isolation. The choice depends on whether consistent latency with MIG or maximum utilization with shared streams is the priority.

13.6.5.2 Model swapping and host memory

When the aggregate size of all models exceeds GPU memory capacity, the serving system must swap models between host dynamic random-access memory (DRAM) and device memory (VRAM) on demand. This introduces a new latency component determined by the PCIe bus bandwidth.


For a 10 GB model on PCIe Gen4 x16 (32 GB/s theoretical bandwidth), loading takes at least 312.5 ms before deserialization, graph setup, or warmup.

To mitigate this, systems use pinned memory (page-locked host memory). GPU direct memory access requires stable physical pages, whereas ordinary host buffers are pageable. Transfers from pageable buffers therefore commonly pass through a temporary pinned staging buffer, adding a CPU copy and latency before the GPU transfer can proceed. Allocating or reusing pinned buffers removes that staging copy, although excessive pinning reduces the memory the operating system can manage freely.

Pinning memory instructs the OS to keep that region permanently in physical RAM. The GPU's DMA engine can then transfer data directly from the pinned region without an extra pageable-to-pinned staging copy. The trade-off is that pinned memory reduces the RAM available for other processes and cannot be reclaimed under memory pressure. For model serving, the transfer-path improvement often justifies pinning model weights and frequently-used input buffers, while leaving less critical memory pageable.

Once a model is loaded, warmed up, and producing predictions consistent with training, request grouping becomes the next optimization lever. Batching directly affects both the throughput and latency terms in our queuing equations.

¹⁸ | **MIG (multi-instance GPU):** Introduced with NVIDIA's A100 (NVIDIA Corporation 2020a) and documented across supported data-center GPUs (NVIDIA 2026e), MIG partitions a single physical GPU into independent instances, each with dedicated compute and memory resources; unlike software sharing (MPS or time-slicing), MIG provides hardware-level isolation between partitions. The trade-off is granularity—partitions must follow fixed profiles, so resources cannot be divided arbitrarily. For multi-model serving, MIG reduces noisy-neighbor risk on shared hardware, while per-model SLO guarantees still depend on scheduler policy, load, and the selected partition profile.



Model loading lands far outside a tight serving SLO.

13.7 Throughput Optimization

Consider a representative ResNet-50 classifier scenario on a V100 GPU at batch size one: the GPU processes one image, then sits idle while the CPU fetches and preprocesses the next—achieving only 15 percent hardware utilization and 200 images per second. The same GPU processing 32 images at once reaches 95 percent utilization and 1,280 images per second, a $6.4\times$ throughput improvement on identical hardware because fixed costs are amortized across requests. The difference is batching, the core lever for improving serving economics. **Batching**¹⁹ differs sharply between training and serving (Crankshaw et al. 2017). *Training* batches maximize throughput by processing hundreds or thousands of samples together with no concern for individual sample latency. *Serving* batches must balance throughput against individual request latency, often processing small batches while ensuring no request waits too long. This adaptive approach is called **dynamic batching** because the system adjusts batch composition in real time based on arriving requests.

Definition 13.5: Dynamic batching

Dynamic batching is the ML serving optimization that trades latency for throughput under stochastic arrival patterns.

1. **Significance:** By buffering requests into a batching window, the scheduler amortizes fixed overheads (L_{lat}) across multiple inputs, pushing the system away from the memory-bound regime (BW) toward the compute-bound regime (R_{peak}).
2. **Distinction:** Unlike fixed-size serving batches, dynamic batching adjusts the realized batch size at inference time according to arriving traffic, a timeout, and a maximum-batch limit.
3. **Common pitfall:** A frequent misconception is that batching “always helps.” In reality, there is a latency-throughput Pareto frontier: if the batching window is too large, the increased queuing delay may violate the system’s SLO before the throughput gains are realized.

13.7.1 Why batching helps

Modern accelerators achieve peak efficiency only at sufficient batch sizes (Shen et al. 2019). A single inference request leaves most compute units idle because GPUs are designed for parallel execution across thousands of threads. Batching amortizes fixed costs across multiple requests and enables parallel execution across the batch dimension.

Two fixed costs dominate at small batch sizes. **Kernel launch overhead**²⁰ is the time for the CPU to prepare and submit work to the GPU. Each layer in a neural network typically requires a separate kernel launch: the CPU must assemble kernel parameters, copy them to GPU-accessible memory, and signal the GPU to begin execution. This overhead is typically $5\text{--}20\ \mu\text{s}$ per kernel, independent of batch size. ResNet-50 has approximately fifty layers, so kernel launch alone adds $250\text{--}1000\ \mu\text{s}$ per inference. At batch size one, this overhead may exceed the actual compute time; at batch size thirty-two, the same overhead is amortized across thirty-two images. Weight loading reads model parameters from GPU memory (VRAM) to the compute units. At batch size one, the GPU reads all weights to process one image; at batch size thirty-two, the same weight read processes thirty-two images, achieving $32\times$ better memory efficiency. Measuring batching efficiency on a concrete model quantifies how these fixed costs amortize in practice.

Table 13.9 shows the throughput-latency trade-off: larger batches can improve hardware efficiency but increase per-request latency. In practice, the optimal batch size depends on both the latency SLO and the arrival rate of requests. The engineering question is quantitative: determine the largest batch size that still meets a given latency SLO. In this scenario, batch size 8 with a 5 ms batching window has worst-case user latency of about 14 ms (5 ms wait plus 9 ms inference), below a 20 ms SLO budget. That earns $4.4\times$ higher service throughput than batch-1 serving on the same hardware, provided sustained load is high enough to fill the batching window. Plotting the same trade-off in figure 13.6 reveals the knee where extra batching stops paying for its latency cost: throughput is already flattening while latency begins to spike, so batching beyond that point trades modest capacity gains for queueing delay.

¹⁹ | **Batch:** From Old French *bache* (a quantity baked at one time), entering computing in the 1950s for jobs processed together without human interaction. The ML serving usage preserves the original trade-off: grouping requests amortizes fixed costs (kernel launch, weight movement) across multiple inputs, but an online request may wait for a batch to form. Training can use batches of hundreds or thousands when memory permits; serving batch sizes range from one to much larger values depending on the model, arrival rate, and latency SLO.

²⁰ | **Kernel (GPU):** CUDA borrowed this term from operating systems circa 2007 because GPU functions represent the computational “core” of parallel algorithms. Unlike OS kernels that run continuously, GPU kernels are discrete units of parallel work launched by the CPU. Each launch carries $5\text{--}20\ \mu\text{s}$ of overhead independent of batch size—negligible for large training batches but dominant at batch-1 serving, where a 50-layer model accumulates $250\text{--}1000\ \mu\text{s}$ of pure launch overhead per inference.




Figure 13.6: The Throughput-Latency Knee: Batch size vs. throughput (blue) and latency (orange). Throughput increases with batch size as hardware utilization improves, but eventually saturates. Latency remains relatively flat until the knee, after which it spikes due to queuing. Values are representative and depend on model/hardware.

The “knee” in figure 13.6 marks the point where the blue throughput curve begins to plateau just as the orange latency curve starts its sharp upward spike. This is an illustrative knee rather than a universal optimum: the selected operating point depends on the latency SLO, arrival process, and cost objective. The numbers are representative rather than tied to a single benchmark.



Napkin Math 13.4: ResNet-50 batching efficiency

Table 13.9 illustrates the throughput-latency trade-off for a ResNet-50/V100 scenario across batch sizes one through thirty-two. The batch sizes, measured inference times, and GPU utilization are the scenario givens; the per-image compute and throughput columns are derived from them.

Math: For the batch-8 row, we derive per-image compute by dividing the batch latency by batch size, $9.1 \text{ ms} \div 8 \text{ images} \approx 1.1 \text{ ms}$ per image, and throughput by dividing batch size by the same latency, $8 \text{ images} \div 9.1 \text{ ms} \approx 879 \text{ img/s}$. We use the same two divisions for every derived row.

Table 13.9: ResNet-50 batching sweep: Illustrative per-image compute, throughput, and GPU utilization across batch sizes one through thirty-two on a V100. Throughput grows 6.4× from batch one to batch thirty-two as GPU utilization climbs from 15 percent to 95 percent, while pure inference time stretches from 5 ms to 25 ms.

Batch Size	Inference Time	Per-Image Compute	Throughput	GPU Util.
1	5 ms	5 ms	200 img/s	15%
4	7.2 ms	1.8 ms	556 img/s	42%
8	9.1 ms	1.1 ms	879 img/s	65%
16	14 ms	0.9 ms	1,143 img/s	85%
32	25 ms	0.8 ms	1,280 img/s	95%

The times shown are pure inference time, excluding queue wait; section 13.7.6 analyzes how user-perceived latency includes batching-window wait.

Systems insight: Batch size thirty-two achieves 6.4× higher throughput than batch size 1. However, user-perceived latency includes both queue wait and inference time. With a 10 ms batching window and 25 ms inference, total latency reaches 35 ms vs. 5 ms at batch size 1.

The efficiency gains from batching come at a cost: requests must wait for the batch to form. This creates a direct tension between throughput optimization (larger batches) and latency minimization (immediate processing). The different batching strategies and their trade-offs govern how engineers tune this balance.

13.7.2 Static vs. dynamic batching

Static batching waits for a fixed batch size before processing, which is simple to implement but fragile under variable traffic: during low traffic, requests wait indefinitely for a full batch, and during high traffic, large batches increase per-request latency. Dynamic batching addresses this failure mode by collecting requests within a bounded time window and processing whatever has arrived when the window closes (Olston et al. 2017; NVIDIA 2024d). The window size becomes the tuning knob: shorter windows reduce latency but sacrifice throughput, longer windows improve throughput but increase latency, and latency-sensitive deployments tune both the time window and maximum batch size against arrival pattern, model shape, and SLO.

13.7.3 Dynamic batching latency-throughput trade-offs

Dynamic batching introduces a quantifiable tension between throughput optimization and latency constraints. Under overload, the mechanism is queue growth rather than slower inference, which enables systematic configuration decisions instead of trial-and-error tuning.

Systems Perspective 13.5: Why latency spikes under load

Recall from section 13.5.1 that Little’s Law ($Q_{\text{req}} = \lambda_{\text{arr}} T_{\text{lat}}$) relates long-run averages only while the queue is stable. As the arrival rate approaches the maximum service rate μ , the stable queue’s depth and latency grow rapidly. If λ_{arr} exceeds capacity for a sustained interval, no finite steady state exists: backlog accumulates until traffic falls, capacity grows, or work is rejected. This is why admission control can preserve latency during overload; scale-out or degraded service can also help when available.

Equation 13.7 decomposes the total user-perceived latency for a batched request into two components:

$$L_{\text{lat}} = L_{\text{lat,wait}} + L_{\text{lat,compute}}(B) \quad (13.7)$$

where $L_{\text{lat,wait}}$ is the time spent waiting in the batching queue (corresponding to $L_{\text{lat,queue}}$ in the overall latency budget) and $L_{\text{lat,compute}}(B)$ is the inference time for batch size B (encompassing $L_{\text{lat,infer}}$ plus portions of $L_{\text{lat,pre}}$ and $L_{\text{lat,post}}$). The batching window T_{window} bounds wait time ($L_{\text{lat,wait}} \leq T_{\text{window}}$), while batch size affects compute time through GPU utilization characteristics.

13.7.3.1 Quantitative analysis of batching

For fixed calendar windows under stationary Poisson arrivals, request arrival times are uniform within each window. A request arriving at time t waits $T_{\text{window}} - t$ for that window to close, so equation 13.8 gives an average wait of half the window:

$$E[L_{\text{lat,wait}}] = \frac{T_{\text{window}}}{2} \quad (13.8)$$

This simple relationship has direct implications. A 20 ms batching window adds 10 ms average wait (up to 20 ms for the first arrival in a window; later arrivals wait less) regardless of batch size achieved. For a 50 ms mean latency SLO with 5 ms inference, the average wait consumes 20 percent of the latency budget before any computation begins; tail SLOs must budget the full window.

13.7.3.2 Batch size distribution

For fixed calendar windows, the number of requests collected during T_{window} follows a Poisson distribution with mean $\lambda_{\text{arr}} T_{\text{window}}$. Equation 13.9 formalizes this relationship:

$$\Pr(\text{batch size} = k) = \frac{(\lambda_{\text{arr}} T_{\text{window}})^k e^{-\lambda_{\text{arr}} T_{\text{window}}}}{k!} \quad (13.9)$$

Here k is a nonnegative integer count of arrivals in the window.

Table 13.10 quantifies this variability, showing how batch size fluctuates for different traffic levels with a fixed 10 ms window:

Table 13.10: Batch Size Variability: At low traffic, fixed calendar windows frequently contain zero requests, so no GPU batch is launched. At moderate traffic, batch sizes fluctuate significantly around the mean. High traffic provides more stable batching, and the probability of batches reaching at least twice the mean size decreases as traffic grows (from 39 percent at 50 QPS to 0.3 percent at 1000 QPS), reflecting the law of large numbers.

Arrival Rate	Mean Batch	Std Dev	Pr(batch = 0)	Pr(batch $\geq 2 \times$ mean)
50 QPS	0.5	0.7	61%	39%
200 QPS	2	1.4	14%	14%
500 QPS	5	2.2	0.7%	3%
1000 QPS	10	3.2	0.005%	0.3%

13.7.3.3 Throughput maximization strategy

Throughput optimization requires separating per-request latency from saturated service capacity. A request waits for batch formation, then pays the service time of the formed batch. Under sustained load, however, batch formation can overlap with the previous batch’s execution, so capacity is governed by the ready-batch service time:

$$\mu_{\text{eff}}(B) \approx \frac{B}{T_{\text{svc}}(B)}, \quad \text{throughput} = \min(\lambda_{\text{arr}}, \mu_{\text{eff}}(B)) \quad (13.10)$$

In equation 13.10, λ_{arr} is the offered arrival rate and $\mu_{\text{eff}}(B)$ is the saturated service capacity for batch size B . The numerator increases linearly with batch size while service time often increases sub-linearly over a useful range because GPU parallelism is better used. The batching window still appears in request latency and in low-traffic regimes, where the expected batch size is limited by arrivals during the window, roughly $\lambda_{\text{arr}} T_{\text{window}}$.

For ResNet-50 on a V100 GPU, service time approximately scales as $T_{\text{svc}}(B) = 5 \text{ ms} + 0.6B$ (5 ms fixed overhead plus 0.6 ms per image in the batch). This linear approximation captures the dominant trend; actual service times may deviate slightly due to memory hierarchy effects. With a $T_{\text{window}} = 10 \text{ ms}$ batching window, table 13.11 extends the pure-inference sweep of table 13.9: latency includes the full-window wait bound, while saturated throughput uses service time alone:

Table 13.11: Batching Throughput Analysis: Scenario analysis for ResNet-50 throughput on V100 with 10 ms batching window. Throughput increases 7.4 $\times$ from batch size one to 32 (179 img/s to 1322 img/s), but total latency more than doubles (15.6 ms to 34.2 ms). The optimal configuration depends on whether the latency SLO or throughput target is the binding constraint.

Batch Size	Service Time	Total Latency	Throughput	Efficiency
1	5.6 ms	15.6 ms	179 img/s	Low
4	7.4 ms	17.4 ms	541 img/s	Moderate
8	9.8 ms	19.8 ms	816 img/s	Good
16	14.6 ms	24.6 ms	1096 img/s	High
32	24.2 ms	34.2 ms	1322 img/s	Maximum

The throughput gains in table 13.11 trace directly back to the fixed-overhead term in the iron law established in section 8.2, where batching amortizes work across requests.



Example 13.7: The iron law of batching efficiency

Context: In serving, batching improves throughput by amortizing fixed per-batch scheduling and kernel-launch work; it can also amortize weight traffic across requests, while queue wait

remains a separate latency cost. Holding the data-movement term outside this illustrative two-term comparison, equation 13.11 isolates compute and fixed overhead:

$$T = \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}} \quad (13.11)$$

Analysis:

- **Case 1 (batch 1):** Overhead (5 ms) $\gg$ Compute (0.6 ms). Efficiency $\approx$ 10 percent. The GPU is mostly waiting.
- **Case 2 (batch 32):** Overhead (5 ms) $\ll$ Compute (19.2 ms). Efficiency $\approx$ 79 percent. The GPU is crunching numbers.

Systems insight: Increase batch size until fixed overhead becomes negligible (< 10 percent of total time) or the latency SLO blocks further waiting. Beyond this point, additional batching yields minimal throughput but imposes a linear queueing penalty.

These three results compose into one working model of dynamic batching: the window sets the average wait at half its length, Poisson arrivals make the realized batch size fluctuate around $\lambda_{\text{arr}} T_{\text{window}}$, and the saturated service capacity $\mu_{\text{eff}}(B)$ climbs with batch size until fixed overhead is amortized away. None of them yet enforces the latency SLO. The passes that follow add that missing constraint, working backward from a hard percentile budget to the largest batch the window may safely form.

13.7.3.4 Latency-constrained optimization

When latency SLOs provide the binding constraint, the optimization problem becomes finding the maximum batch size that meets the SLO. For a latency target $L_{\text{lat,target}}$ and average wait time $T_{\text{window}}/2$, equation 13.12 defines the maximum allowable batch size using a first-order *average* latency approximation:

$$B_{\text{max}} = \max \left\{ B : \frac{T_{\text{window}}}{2} + T_{\text{svc}}(B) \leq L_{\text{lat,target}} \right\} \quad (13.12)$$

In this illustrative ResNet-50 scenario with a 50 ms p95 latency SLO, comparing a conservative batching window against an aggressive one shows how the assumed configurations trade wait time, inference budget, and saturated capacity. Table 13.12 lays the two configurations side by side.

Table 13.12: Illustrative batching window trade-off: How two assumed batching configurations trade average wait, inference budget, batch-size cap, and saturated capacity.

Metric	Conservative ($T_{\text{window}} = 5$ ms)	Aggressive ($T_{\text{window}} = 25$ ms)
Average wait	2.5 ms (max wait = 5 ms for the first request in a window)	12.5 ms
Latency budget for inference	47.5 ms (mean-latency planning; tail SLOs should budget the full window)	37.5 ms
Batch size cap	32 images	48
Assumed saturated capacity	$\sim 1,140$ img/s	$\sim 1,280$ img/s

The aggressive configuration assumes 12.3 percent higher saturated capacity but increases average batching wait by 10 ms and the maximum window-wait contribution by 20 ms. The latter is not a p99-latency estimate.

13.7.3.5 SLO violation analysis

Batch size variability can cause SLO violations even when mean latency appears safe. A request in a fixed window sees itself plus the other arrivals, so $B_{\text{req}} = 1 + \text{Poisson}(\lambda_{\text{arr}} T_{\text{window}})$. Combining the full window with the request-seen p99 batch size gives the conservative envelope in equation 13.13:

$$L_{\text{lat,envelope}} = T_{\text{window}} + T_{\text{svc}}(B_{\text{req,p99}}) \quad (13.13)$$

For $\lambda_{\text{arr}} = 500$ QPS and $T_{\text{window}} = 10$ ms, the mean request-seen batch size is 6 and its p99 is 12. The mean latency is 13.6 ms; the conservative envelope is 22.2 ms.

The envelope is 1.63× the mean. It is not the request-latency p99 because full-window wait and p99 batch size need not occur for the same request.

🔍 Systems Perspective 13.6: The latency-throughput trade-off

A single “inference speed” number is undefined until the batch size is named, because batch size selects which regime, and which bottleneck, the system operates in.

- **Batch-1 regime:** Latency-oriented. Launch overhead or memory traffic often limits the request path. This regime governs real-time interaction such as typing helpers and robotics.
- **Batch-N regime:** Throughput-oriented. Larger batches amortize overhead and weight traffic and may shift the bottleneck toward compute. This regime governs offline processing and high-traffic services.

The two regimes optimize opposite quantities, so a model that is “fast” at batch 1 may be far from peak throughput, and vice versa. Any latency or throughput figure must therefore specify whether it was measured at single-stream latency (batch 1) or maximum throughput (batch N).

13.7.3.6 Adaptive batching windows

The same batch-size dependence drives how the serving system shapes its batches in the first place. Fixed batching windows waste latency budget during high traffic when large batches form quickly. Listing 13.2 demonstrates how adaptive strategies adjust the window based on queue depth.

In an illustrative scenario with traffic varying between 200–1000 QPS, a fixed 10 ms window produces 15 ms average latency at 650 img/s, while the adaptive heuristic produces 11 ms at 680 img/s. The interplay between window size and batch limits creates a space of possible configurations, each representing a different balance between throughput and latency.

Listing 13.2: Adaptive Batching Window: Heuristically adjusts batch timeout based on queue depth, arrival rate, and the remaining latency budget.

```
def adaptive_batching_window(
    queue_depth, arrival_rate, slo_ms, service_ms, fixed_overhead_ms
):
    """Compute a heuristic batching window.

    Based on current system state.
    """
    target_batch_size = 16 # Deployment-tuned target batch

    # Fast path: batch ready, close immediately to minimize latency
    if queue_depth >= target_batch_size:
        return 0

    # Compute maximum allowable wait from the remaining p99 budget.
    max_wait_ms = max(0, slo_ms - service_ms - fixed_overhead_ms)

    # Estimate time to accumulate target batch at current arrival
    # rate.
    # arrival_rate is requests/second, so convert seconds to
    # milliseconds.
    if arrival_rate > 0:
        requests_needed = target_batch_size - queue_depth
        estimated_wait_ms = requests_needed / arrival_rate * 1000.0
        # Return minimum of estimated wait and SLO-constrained maximum
        return min(estimated_wait_ms, max_wait_ms)

    return max_wait_ms # Low traffic: use remaining budget to accumulate batch
```


The nondominated batching configurations form a Pareto frontier where improving throughput requires accepting higher latency. Table 13.13 shows five illustrative configurations:

Table 13.13: Illustrative Batching Pareto Frontier: The five assumed configurations trade tail latency for throughput. From the first to the last, assumed throughput rises by 52 percent and assumed p99 latency rises to 5.4× its initial value.

Window (ms)	Max Batch	Avg Latency	p99 Latency	Throughput	Configuration
2	16	8 ms	18 ms	890 img/s	Ultra-low latency
5	32	10 ms	22 ms	1,140 img/s	Balanced
10	32	15 ms	35 ms	1,240 img/s	Moderate latency
20	64	23 ms	52 ms	1,310 img/s	Throughput-optimized
50	128	38 ms	98 ms	1,350 img/s	Maximum throughput

13.7.3.7 Practical configuration guidelines

The Pareto frontier in table 13.13 illustrates why configuration must work backward from the latency budget: past the knee, widening the window buys diminishing throughput for sharply rising tail latency. In the worked scenario, reserving 30 percent of the SLO for batching wait gives $T_{\max} = 0.3 \times L_{\text{lat,SLO}}$; another workload may allocate a different fraction after accounting for network, preprocessing, inference, and postprocessing. The traffic estimate should include peaks and burst duration rather than rely only on the mean. GPU memory supplies a separate ceiling on batch size through activations and runtime workspace. Finally, monitoring the realized batch-size and wait-time distributions reveals whether the assumed arrival model holds and whether a fixed or adaptive window is appropriate.

For an illustrative ResNet-50 serving system with a 50 ms SLO and 500 QPS traffic, the calculation turns the SLO and arrival-rate assumptions into two deployable knobs: the batching window and maximum batch size. Table 13.14 summarizes the resulting configuration and modeled operating point.

Table 13.14: Practical Batching Configuration: A 30 percent allocation sets a 15 ms maximum window; this illustrative scenario selects a 12 ms window below that bound and assumes a batch-32 cap. The reported latency is a conservative envelope, not the request-latency p99.

Quantity	Value	Engineering role
Latency budget for batching	15 ms	Portion of the SLO available for queueing delay.
Maximum window	15 ms	Upper bound implied by the latency budget.
Expected request-seen batch	7	Average batch seen by an arriving request.
Request-seen p99 batch	13	Tail batch size under Poisson arrivals.
Assumed memory-limited batch	32	Scenario cap for accelerator memory.
Selected configuration	$T_{\text{window}} = 12 \text{ ms}$, $B_{\max} = 32$	Practical knob setting for deployment.
Conservative tail envelope	24.8 ms	Full window plus request-seen p99 batch service time.
Illustrative saturated capacity	1,176.9 img/s	Estimate using the assumed 0.89 efficiency factor.
Served throughput	500 img/s	Arrival-limited load handled by the configuration.

13.7.4 Continuous batching

Autoregressive models like language models generate outputs token by token: each new token depends on all previously generated tokens, so generation is inherently sequential. The dynamic batching examined earlier in this section assumes fixed-length outputs. LLMs violate this assumption: if one sequence in a batch of eight finishes after ten tokens while others need 100 tokens, 90 percent of the compute for that sequence slot is wasted (Yu et al. 2022).

Continuous batching²¹ (also called iteration-level batching) addresses this waste by allowing new requests to join a batch between generation steps and completed sequences to exit (Yu et al. 2022; Kwon et al. 2023). The system manages batch composition dynamically at each decoding iteration rather than forming static batches that persist for the entire generation process.

The mechanism works as follows: when a sequence generates its end-of-sequence token, its slot becomes immediately available. A waiting request can fill that slot for the next iteration rather than

²¹ | **Continuous batching:** Also called “iteration-level batching” (Yu et al. 2022) and, in NVIDIA TensorRT-LLM, “in-flight batching” (NVIDIA 2026g); the distinction is scheduling granularity. Traditional batching commits to a fixed batch for an entire generation sequence (potentially hundreds of iterations), while continuous batching reschedules at every token-generation step—analogous to preemptive OS process scheduling vs. run-to-completion. This finer granularity reduces the waste from variable-length sequences, where a batch slot occupied by a completed sequence sits idle until all other sequences finish.

waiting for the entire batch to complete. Similarly, the system can add new requests to available slots without interrupting ongoing generation. This dynamic approach maintains high GPU utilization even when sequence lengths vary widely.

Systems implementing continuous batching, such as vLLM²² and TensorRT-LLM, improve throughput by keeping decode slots occupied as sequences enter and exit (Kwon et al. 2023; NVIDIA 2026g). Sarathi-Serve refines this scheduler with chunked prefill and stall-free batching to reduce interference between prompt processing and token decoding (Agrawal et al. 2024). The improvement comes from two sources: reducing wasted compute on completed sequences and reducing average wait time for new requests. For production language model serving where response lengths vary from single tokens to thousands, continuous batching has become a central technique for cost-effective deployment.


Memory management adds complexity to continuous batching. As sequences enter and exit the batch, the key-value cache that stores attention context must be dynamically allocated and freed. Consider what happens when sequences of varying lengths share GPU memory: a 100-token sequence completes and releases its cache, but a new 150-token sequence cannot use that space because it needs a larger contiguous block. Over time, small unusable gaps accumulate between allocated regions, eventually preventing new sequences from starting even when total free memory appears sufficient. This memory fragmentation can waste 40 to 50 percent of available memory in naive implementations, severely limiting the concurrent batch size that determines throughput.

13.7.4.1 PagedAttention

PagedAttention,²³ introduced in vLLM, addresses this fragmentation problem by applying operating system virtual memory concepts to GPU memory (Kwon et al. 2023). In static contiguous allocation, reserving maximum sequence length memory up front causes internal waste, while freed regions may be too small or poorly placed to satisfy a new contiguous reservation. Instead of allocating one contiguous region per sequence, PagedAttention divides the KV cache into fixed-size blocks mapped through a block table. A sequence's cache can therefore occupy noncontiguous physical blocks. When a sequence completes, its blocks return to a free pool for reuse. The vLLM paper reports KV-cache waste below 4 percent and higher throughput than the compared systems; exact gains depend on workload and implementation.

The batching and memory techniques covered here establish the foundation for LLM serving, but several advanced topics warrant additional study.

²² | **vLLM (virtual LLM)**: An open-source serving system that enables continuous batching via its PagedAttention algorithm (Kwon et al. 2023). Inspired by OS virtual memory, this technique reduces the severe KV-cache fragmentation and reservation waste that constrains static batching. By keeping KV-cache waste low, vLLM can serve larger effective batches on the same hardware.



PagedAttention recovers the KV-cache waste of contiguous allocation.

²³ | **PagedAttention**: The name directly references OS virtual memory paging, first implemented on the Atlas computer at Manchester (1962) to solve the same class of allocation problem—programs needed more memory than physically available, and contiguous allocation wasted space (Kilburn et al. 1962). Introduced at SOSP 2023, PagedAttention applies this six-decade-old abstraction to GPU KV-cache memory: before it, LLM serving systems wasted 60–80 percent of KV cache memory due to fragmentation and over-reservation. PagedAttention reduces waste to under 4 percent, enabling 2–4× higher throughput on the same hardware (Kwon et al. 2023).

🔍 Systems Perspective 13.7: LLM serving: Beyond the fundamentals

Language model serving introduces challenges beyond the batching and memory principles established here. The key-value cache that stores attention context scales with sequence length and batch size, often exceeding the model weights themselves in memory consumption. Techniques like speculative decoding use small draft models to propose multiple tokens that the target model verifies in parallel; evaluated T5-XXL workloads achieved roughly 2–3× speedups (Leviathan et al. 2023), but gains vary. Weight-only quantization (such as INT4 weights with FP16 activations) is especially relevant for memory-bandwidth-bound LLM inference (Lin et al. 2024).

These LLM-specific optimizations build directly on the foundations this chapter establishes: queuing theory governs request scheduling, batching trade-offs determine throughput-latency curves, and precision selection follows the same accuracy-efficiency principles. The same serving fundamentals remain applicable, while LLM serving adds domain-specific techniques atop them. Advanced treatments provide detailed coverage of KV cache optimization, including techniques for multi-tenant serving, where one fleet shares capacity across users, and distributed inference, where one request may be split across machines.

Continuous batching is widely used for high-throughput LLM serving, yet not all deployment scenarios benefit from batching. The techniques examined so far, from dynamic batching windows to PagedAttention, target shared server workloads. They introduce scheduler and memory-

management complexity that may not be justified in every deployment. The practical question is *when* batching hurts rather than helps.

Some scenarios favor single-request processing. Under an ultra-low p99 target, even a short batching timeout may consume too much of the budget. Highly variable request sizes create padding overhead because shorter inputs may be padded to match the longest input in a batch. Memory can also bind when weights already consume most accelerator memory, since batched activations and cache state can trigger out-of-memory errors.

13.7.5 Session affinity constraints

When requests from the same user or session should route to the same replica, batching becomes constrained. Session affinity, also called sticky sessions, matters for three main reasons.

The most impactful case is KV-cache reuse in conversational AI, where the key-value cache from previous turns can materially speed up multi-turn conversations. Routing a follow-up request to a different replica forfeits this cached context, forcing the system to recompute or reload prefix state for long conversations.

A second driver is user-specific models: some systems serve personalized models or adapters per user, and routing requests to the replica that has already loaded that user's adapter avoids repeated loading overhead. Similarly, stateful preprocessing that maintains tokenizer caches or session-specific normalization requires rebuilding state when requests route to a different replica.

The tension with batching is clear since strict affinity constrains which requests can be batched together, potentially reducing batch sizes and GPU utilization. Production systems often implement soft affinity where requests prefer their assigned replica but can overflow to others when that replica is overloaded. This preserves most affinity benefits while maintaining load balance.

13.7.6 Traffic patterns and batching strategy

The batching strategy depends critically on how requests arrive. Different deployment contexts exhibit different arrival patterns, and the MLPerf inference benchmark supplies four standardized scenarios that approximate several common cases, as section 12.8.4.2 explains in detail.

13.7.6.1 Server traffic (Poisson arrivals)

The MLPerf Server scenario models cloud/API-like inference traffic with Poisson arrivals (Reddi et al. 2019).²⁴ Under that model, counts in disjoint intervals are independent and the average arrival rate is constant. Equation 13.14 expresses the expected count in a fixed calendar window with rate λ_{arr} and width T_{window} :

$$E[\text{batch size}] = \lambda_{\text{arr}} \cdot T_{\text{window}} \quad (13.14)$$


For fixed calendar windows, variance equals the mean, so batch sizes fluctuate significantly at moderate traffic. With $\lambda_{\text{arr}} = 200$ requests/second and $T_{\text{window}} = 10$ ms, the expected count is two, but roughly 14 percent of windows have zero requests (and launch no GPU work) while others may have four or more.

A useful heuristic for the batching window balances waiting cost against throughput benefit. Equation 13.15 expresses one such rule:

$$T_{\text{window}} \approx \min \left(L_{\text{lat,SLO}} - T_{\text{svc}}, \sqrt{\frac{T_{\text{svc}}}{\lambda_{\text{arr}}}} \right) \quad (13.15)$$

where $L_{\text{lat,SLO}}$ is the latency SLO, T_{svc} is the service time (in seconds), and λ_{arr} is the arrival rate (in requests per second), making the second term dimensionally consistent in seconds. The square-root form is a local cost-model heuristic: it balances a fixed per-batch benefit against a waiting cost that grows with the arrival interval. It is not a closed-form optimum for ML serving specifically; production systems calibrate the window empirically against observed traffic. A counterintuitive result emerges from equation 13.15: as traffic increases, the optimal window decreases while achieved batch sizes still grow. Table 13.15 demonstrates this phenomenon across four traffic levels.

²⁴ | **Poisson process:** Named after French mathematician Simeon Denis Poisson (1781–1840), this stochastic model describes events occurring continuously and independently at a constant average rate. In fixed calendar windows, variance equals the mean, so with $\lambda_{\text{arr}} = 200$ req/s and a 10 ms window, the expected count is two and roughly 14 percent of windows are empty, though an efficient scheduler launches no GPU work for them. A request-triggered timer instead begins with one request and counts only additional arrivals during the window.



As QPS rises, the batching window shrinks while batch size grows.

Table 13.15: Traffic-Adaptive Batching: Higher traffic enables shorter windows while still achieving larger average batches. Values are computed from equation 13.15 with a 50 ms SLO and a 25 ms service-time assumption, so the latency column is the approximate service-plus-window budget rather than a measured production p99.

Arrival Rate	Optimal Window	Avg Batch Size	Approx. Latency
100 QPS	15.8 ms	1.6	40.8 ms
500 QPS	7.1 ms	3.5	32.1 ms
1,000 QPS	5 ms	5	30 ms
5,000 QPS	2.24 ms	11.2	27.2 ms

13.7.6.2 Streaming traffic (correlated arrivals)

Autonomous vehicles, video analytics, and robotics systems receive inputs from multiple synchronized sensors. A six-camera example makes the synchronization deadline concrete.



Example 13.8: Multi-camera autonomous vehicle serving

Scenario: An autonomous vehicle vision system processes 6-camera video streams under a strict 33 ms real-time frame deadline.

Diagnosis: Sensor arrival jitter causes variable multi-camera batch assembly delays, risking frame deadline misses. Implementing a 12 ms synchronization budget and fallback rules (reusing previous frames on timeout) bounds latency.

Systems lesson: Synchronized multi-sensor streaming workloads enforce deterministic real-time deadlines. Serving pipelines must implement explicit synchronization timeout policies to prevent sensor jitter from breaching latency budgets.

Table 13.16: Multi-camera frame timeline: Event-by-event timeline for one synchronized frame set across six cameras at 30 FPS. Frames arrive over a 7 ms spread, inference cannot start until the last one lands at 15 ms, and the result reaches the planning module at 32 ms, leaving 1 ms of slack inside the 33 ms deadline.

Time	Event
$T = 0$ ms	Cameras begin capturing frame N
$T = 8$ ms	Camera 1 frame arrives
$T = 10$ ms	Cameras 2–5 frames arrive
$T = 15$ ms	Camera 6 arrives (jitter)
$T = 15$ ms	Batch inference begins (6 images)
$T = 25$ ms	Inference complete
$T = 32$ ms	Result ready for planning module

13.7.6.3 Single-user traffic (sequential arrivals)

Streaming traffic correlates arrivals through sensor synchronization, making batch size and deadline externally fixed. At the opposite end of the spectrum, some interactive mobile and embedded applications have little opportunity to batch independent requests. The optimization target shifts from a synchronization budget under a hard frame deadline to per-request latency and energy consumption under a thermal power envelope.

Many interactive mobile and embedded applications process one request stream at a time; the MLPerf SingleStream scenario captures this sequential-serving shape. For the ResNet-50 phone scenario used here, the dominant costs shift from batch formation to per-request latency and energy.

13.7.6.4 Mobile serving constraints

Cloud services often emphasize fleet cost and latency; mobile serving adds three closely coupled constraints. The first is an energy budget that throughput alone does not capture, because each inference consumes battery. In the modeled pipeline, 2.98 mJ at 22 FPS draws about 66 mW for the inference path alone, before camera, display, ISP, and OS overhead add to that total in a full photo app. Thermal throttling can compound this limit when sustained system power reaches a

device-specific thermal boundary. Memory constraints close the set because the model shares RAM with the application and operating system. Reload time depends on device storage, runtime, and model state; quantization and memory-mapped loading (section 13.6.3) can reduce footprint and startup work but do not eliminate those dependencies.



Example 13.9: ResNet-50: Mobile serving

Table 13.17 decomposes per-phase latency and energy for a single-user mobile vision inference:

Table 13.17: Mobile ResNet-50 pipeline: Illustrative per-phase latency and energy for a single-user mobile vision inference, showing that JPEG decode on the CPU dominates the energy budget even though the NPU inference stage carries the model's compute. Optimization targets shift from throughput to energy-per-inference at the edge.

Phase	Duration	Energy	Notes
Camera buffer read	8 ms	0.08 mJ	System API
JPEG decode (CPU)	15 ms	1.5 mJ	Single-threaded
Resize + Normalize	5 ms	0.4 mJ	CPU preprocessing
NPU inference	12 ms	0.8 mJ	82% utilization
Postprocess + UI	5 ms	0.2 mJ	Result rendering
Total	45 ms	2.98 mJ	22 FPS sustained

The mobile serving node is governed by four metrics:

- **Energy per inference:** 2.98 mJ enables ~12.1M inferences per 10 Wh battery (typical smartphone)
- **Modeled inference power:** 2.98 mJ / 45 ms = 66 mW, excluding the rest of the device and therefore not a guarantee of indefinite operation
- **NPU vs. CPU trade-off:** CPU fallback replaces the 12 ms, 0.8 mJ NPU inference stage with a 45 ms, 4.2 mJ CPU stage; the full pipeline would rise from 45 ms and 2.98 mJ to about 78 ms and 6.4 mJ before additional system overhead.
- **Memory footprint:** 150 MB peak (model + activations), competing with app memory

Systems insight: In this illustrative pipeline, the mobile NPU reaches 82 percent utilization at batch size one while the data-center example reaches 15 percent. The contrast is workload- and hardware-specific, but it shows why a mobile accelerator designed for single-stream execution can operate efficiently without the batches used to amortize a large server GPU.

These constraints make mobile serving optimization qualitatively different from cloud optimization. The goal is not maximum throughput but sustainable performance, maintaining acceptable latency without thermal throttling or excessive battery drain. Strict latency budgets make synchronized multi-camera frame processing challenging (table 13.16). In a 30 FPS autonomous driving pipeline (33 ms total budget), frame capture and arrival jitter consume 15 ms before inference can begin. Batching 6 camera streams into a single GPU invocation completes in 10 ms (at $T = 25$ ms), leaving 7 ms for postprocessing, planning, and safety checks. If frame jitter pushes the start past 15 ms, the pipeline misses its deadline.

13.7.6.5 Traffic pattern summary

Traffic-adaptive batching adjusts the batching window as queue depth and request rate change. Table 13.18 relates the four MLPerf scenarios to representative deployment contexts and common batching strategies.

These scenarios are benchmark abstractions rather than deployment prescriptions. The operative batching policy must still follow the service's current concurrency, deadline, and power constraints.

Table 13.18: Traffic Patterns and Batching Strategies: The four MLPerf inference scenarios approximate distinct deployment contexts. Server traffic (cloud APIs) uses dynamic batching with a timeout; MultiStream approximates synchronized sensor inputs; SingleStream processes requests individually; Offline uses large batches subject to memory capacity to maximize throughput.

Scenario	Context	Strategy	Focus
Server	Cloud APIs, web services	Dynamic batching with timeout	Window tuning, utilization-latency curve
MultiStream	Autonomous driving, video analytics	Synchronized sensor fusion	Jitter handling, deadline guarantees
SingleStream	Mobile apps, embedded devices	No batching ($B = 1$)	Preprocessing, power efficiency
Offline	Batch processing, data pipelines	Maximum batch size	Throughput, hardware utilization

The MLPerf Server scenario captures cloud API traffic, MultiStream captures synchronized sensor workloads, and Offline inference captures batch processing where throughput dominates latency.

The batching strategies examined so far share a critical assumption: each request produces a single, fixed-size output—one classification label, one bounding box, one embedding vector. This assumption governs the queuing math, the Pareto frontier analysis, and the traffic-adaptive window tuning. The fastest-growing category of serving workloads, however, violates this assumption entirely. Large language models generate outputs token by token, with each token depending on every previous one. A single request may produce hundreds or thousands of tokens over seconds of elapsed time, yet must feel responsive from the first token onward. This fundamental shift from *fixed-output* to *variable-length, streaming-output* serving builds directly on the continuous batching and KV-cache paging already established for autoregressive generation. What it adds are phase-split metrics for prefill and decode, decoding strategies that trade output quality against per-token cost, and memory tactics such as prefix reuse and offloading that exploit shared context.



Checkpoint 13.3: Batching and traffic patterns

Batching is the primary lever for serving economics, but the optimal strategy depends on context.

- ☐ **Throughput-latency trade-off:** can you explain why batch size 32 achieves 6× higher throughput than batch size one, yet a production system with a 20 ms SLO might still choose batch size eight?
- ☐ **Dynamic vs. static batching:** can you describe why static batching (waiting for a full batch) fails under variable traffic, and how dynamic batching with a time window solves this?
- ☐ **Traffic pattern matching:** given a deployment scenario (for example, cloud API, autonomous vehicle, mobile app), can you select the appropriate MLPerf scenario and explain why that batching strategy fits?
- ☐ **Adaptive windows:** can you explain why the optimal batching window *decreases* as traffic *increases*, even though batch sizes grow?

²⁵ | **Autoregressive:** From Greek *auto-* (self) and Latin *regressus* (a going back)—the output “regresses” on itself; George Udny Yule introduced autoregressive models in 1927 for analyzing sunspot cycles. In language modeling, each output token conditions on all previously generated tokens, creating a serial dependency that prevents the parallelism exploited during training. Decode is often weight-bandwidth-bound at small batches; larger batches amortize the weight stream and can shift the bottleneck toward compute or KV-cache traffic (Pope et al. 2023).

13.8 LLM Serving

Large language models make three properties far more prominent than in fixed-output serving: *autoregressive generation*²⁵ (each token depends on all previous tokens, making output inherently sequential), *variable-length output* (response length is unknown at request time, complicating fixed-batch assumptions), and *stateful memory* (the key-value cache grows with each generated token, creating dynamic memory pressure). Together, these properties create a qualitatively different serving challenge. The p50, p95, and p99 metrics used for classification serving still matter, but they apply to different *phases* of the request—the initial prompt processing and the subsequent token generation. The same principles of queuing theory, batching trade-offs, and latency budgets remain applicable; LLM serving adds domain-specific techniques atop them.

13.8.1 Performance metrics: TTFT and TPOT

Generative models produce a stream of tokens rather than a single output tensor. This streaming nature requires dedicated LLM performance metrics that reflect the transition from *prefill* (processing the prompt and populating its KV cache) to *decode* (generating output). Two key measures are **time to first token (TTFT)** and **time per output token (TPOT)**, which capture responsiveness and average generation pace. Per-gap **Inter-Token Latency (ITL)** measures the distribution of pauses between individual streamed tokens; TPOT instead averages the post-first-token interval across an output.

Definition 13.6: LLM performance metrics

LLM performance metrics separate initial responsiveness from the cadence of streaming autoregressive generation.

1. **Significance:** They decompose user-perceived latency into time to first token (TTFT) (whose prefill work is often compute-bound) and time per output token (TPOT) (whose decode work is often bandwidth-bound at small batches).
2. **Distinction:** Unlike fixed-output metrics (for example, end-to-end latency), LLM metrics measure the fluidity of generation, acknowledging that the user experience depends on the *rhythm* of token arrival.
3. **Common pitfall:** A frequent misconception is that a “fast model” has a low TTFT. In reality, a model can have a fast TTFT but a sluggish TPOT (if the memory wall (BW) is the bottleneck), leading to a frustrating user experience where the answer starts quickly but “stutters” thereafter.


These metrics capture distinct aspects of the user experience, so interactive services often set separate targets for TTFT, TPOT or ITL, and aggregate throughput.

Systems Perspective 13.8: LLM serving latency targets

Interactive LLM services usually need separate SLOs for responsiveness, generation fluidity, and fleet throughput. The following values are illustrative targets rather than universal requirements:

- **TTFT:** < 500 ms (for a 1000-token prompt)
- **TPOT:** < 50 ms (equivalent to ~20 tokens/s, faster than human reading speed)
- **Throughput:** > 1000 tokens/s aggregate across active serving replicas

A single “latency” number therefore hides the prefill/decode split. TTFT is the blank-screen budget, TPOT is the reading-flow budget, and aggregate service throughput, measured as tokens/s summed across active serving replicas, determines whether those targets hold under shared load.



13.8.2 Decoding strategies

Meeting these TPOT targets depends on more than memory bandwidth alone: the algorithm used to select each token also affects per-token latency and output quality. Generative models require decoding strategies that trade off quality, diversity, and latency. The choice of decoding strategy dramatically affects both output quality and computational cost.

The simplest approach, greedy decoding, selects the highest-probability token at each step at the cost of one model pass per token. It is fast but can produce repetitive outputs because it cannot recover from early choices. Beam search maintains multiple candidate sequences and selects a high-scoring complete sequence, increasing decoder work and state relative to greedy decoding. Sampling with temperature, top-*k*, and top-*p* (also called nucleus sampling) injects controlled randomness for diversity (Holtzman et al. 2020). The serving cost of any strategy depends on its implementation and on the output-length distribution that continuous batching must absorb.

Interactive LLM services commonly return tokens as they are produced rather than waiting for complete generation. This streaming response transforms the user experience: a two-second total generation feels responsive when tokens stream continuously, but feels broken when users stare at a blank screen for two seconds. Streaming requires infrastructure support for incremental responses and client-side rendering. The latency profile shifts accordingly: TTFT determines when output starts appearing (responsiveness), while TPOT determines the perceived generation speed (fluidity). Once generation is streamed token by token, the serving unit changes from one fixed-output prediction to a stateful sequence whose memory footprint grows on every step.

13.8.3 Memory and KV cache

26 | KV cache (key-value cache): To avoid redundant work, the system caches the Key and Value vectors from previous tokens, which remain valid throughout generation; this design choice is the direct cause of the dynamic memory growth described; the cache's size grows linearly with every generated token, making memory management a primary constraint. For the 70-billion-parameter-class grouped-query-attention sizing example in this calculation, the FP16 KV cache is about 0.33 MB per token per sequence; grouped-query attention (GQA), used in Llama-family models such as Llama 3, shares key/value heads across multiple query heads, reducing the cache relative to full multi-head attention (Dubey et al. 2024). A batch of 32 requests at an 8,000-token context therefore requires roughly 80 GB just for KV cache, several times larger still without grouped-query or multi-query attention.

Generative inference requires managing the KV Cache²⁶, a stateful memory structure that grows with sequence length. Prefill creates cache entries for prompt tokens, and each decoded token adds another entry. Across concurrent variable-length sequences, this dynamic state can dominate available accelerator memory and create allocation waste if it is not managed explicitly.

13.8.3.1 Prefix caching and memory offloading

The continuous batching and PagedAttention techniques covered in section 13.7.4 address request scheduling and cache paging; the remaining memory pressure can be further mitigated through architectural strategies that exploit request patterns. **Prefix caching** stores the KV states of common instruction prefixes (such as a 2,000-token system prompt or a shared retrieval-augmented generation (RAG) context), allowing independent requests to reuse the prefill result for an identical prefix. For N requests sharing a prefix of S_{prefix} tokens, the avoided prefill work is roughly $(N - 1)S_{\text{prefix}}$ token steps, subject to cache hits, eviction, and implementation overhead. For multi-turn conversations, the same principle lets a subsequent turn process only the new suffix after reusing compatible prior state.

Napkin Math 13.5: The energy cost of a chat

Problem: How much energy does an H100-backed chat service spend per generated token and for a response with 500 tokens?

As LLMs scale, joules per token becomes a first-class operational metric alongside latency. Assuming the H100 draws its 700 W thermal design power (TDP), this scenario estimate follows from throughput and power (Choquette 2023):

1. **Throughput:** 114 concurrent requests $\times$ 8 tokens/s per request $\approx$ 912 tokens/s.
2. **Power:** 700 W (GPU) + 300 W (Host/Overhead) = 1000 W.
3. **Energy per token:** 1000 W / 912 tokens/s $\approx$ 1.0965 J/token

Systems insight: Under these assumptions, a response of 500 tokens consumes $\approx$ 548.2 J.

- For comparison, charging a smartphone consumes $\approx$ 40000 J.
- Boiling a cup of water consumes $\approx$ 100000 J.

The primary way to reduce J/token is to increase hardware utilization and eliminate redundant compute. If the GPU sits at 10 percent utilization due to poor batching, it still draws $\sim$ 300 W and the host adds 300 W, causing energy per token to rise to 6.6 J/token (approximately 6 $\times$ the baseline). Architectural optimizations like prefix caching also skip the energy-intensive prefill phase for shared context, directly reducing the energy footprint of retrieval-augmented generation (RAG) and chat applications. The serving lesson is that efficiency is not only a latency or cost metric; it also determines how much energy each useful token consumes.

When the aggregate KV cache exceeds GPU VRAM, systems can employ **KV cache offloading**. This strategy spills inactive or low-priority context windows to host CPU RAM or NVMe SSD, freeing VRAM for active generation. The reload cost is bounded below by bytes moved divided by PCIe or NVMe bandwidth, before software overhead and queueing are added. Offloading therefore prevents OOM failures and enables larger context windows, but it also creates affinity, invalidation, and hot-decode latency risks that must be budgeted explicitly. Advanced techniques including

speculative decoding²⁷ and distributed parallelism, where one request is split across multiple devices or machines, are covered in specialized treatments of large-scale systems.



Checkpoint 13.4: LLM serving fundamentals

LLM serving introduces constraints absent from traditional model serving.

- ❑ **TTFT vs. TPOT:** can you explain why these metrics capture different user experiences, why prefill is often compute-bound while low-batch decode is often bandwidth-bound, and convert a 50 ms TPOT target into tokens per second?
- ❑ **Memory wall:** can you explain why adding compute cores yields little latency improvement when token generation is bandwidth-bound, and why faster memory, smaller models, or batching can change the bound? (The Llama-3 case study in section 13.11.4 quantifies this relationship.)
- ❑ **Continuous batching:** can you explain why traditional static batching wastes compute when sequence lengths vary, and how iteration-level batching solves this?
- ❑ **PagedAttention:** can you explain the memory fragmentation problem in KV-cache management and how fixed-size, noncontiguous blocks reduce reservation waste?
- ❑ **Prefix caching:** can you explain how caching the KV states of common instruction prefixes reduces redundant computation and speeds up RAG or multi-turn applications?

The chat-energy calculation makes utilization part of the LLM serving budget: the cost of each generated token depends on how requests share the machine. Unlike fixed-output inference, LLM energy grows with response length. Each decode step streams weights once for the active batch, so batching amortizes that traffic across sequences. Carbon accounting additionally requires an emissions factor and a defined boundary.

Energy efficiency depends on the same batching, memory, and prefix-cache mechanisms that govern LLM latency, so the useful summary is a constraint checklist rather than a single scalar metric.

13.9 Inference Runtime Selection

The batching strategies and LLM-specific techniques determine *how* requests are grouped and processed. These strategies assume an underlying execution engine that actually runs the model computations. The token generation dynamics introduced in section 13.8 and the latency budgets established in section 13.4.1 are achievable only if the runtime maps operations efficiently to hardware. The **inference runtime**, the software layer that orchestrates tensor operations and manages hardware resources, can materially change performance for the same model. Runtime work therefore has two phases: selection chooses the execution engine, and configuration tunes that engine for the target model, hardware, input shapes, and latency distribution.

13.9.1 Runtime ecosystem and configuration

Selection should start with the binding constraint rather than the framework used during training. When deployment speed and framework compatibility dominate, PyTorch and TensorFlow models can often serve through their native runtimes without a separate interchange export. This preserves support for more of the training framework's operators and control flow, although custom operations, device support, and production packaging can still limit deployability. Framework runtimes also carry generality that a target-specific inference engine may trade for tighter optimization.

`torch.export` and TensorFlow SavedModel provide serialized graphs for ahead-of-time transformation and graph optimization while maintaining framework compatibility. TorchScript remains a legacy format for existing deployments.

13.9.1.1 General-purpose optimization

When portability across hardware is the binding constraint, ONNX Runtime²⁸ provides a hardware-agnostic optimization layer (Microsoft 2024b). Models export to ONNX format, then ONNX Runtime

²⁷ | **Speculative decoding:** A small “draft” model generates k candidate tokens autoregressively; the large target model then verifies the proposed block in parallel (Leviathan et al. 2023). When the draft model's proposals are accepted at rate α , effective throughput can scale with the number of accepted tokens per verification step. This breaks the serial autoregressive bottleneck at the runtime layer, not the architecture layer.

²⁸ | **ONNX Runtime:** Microsoft's inference engine acts as a hardware abstraction layer: the same ONNX model can run on CPUs, NVIDIA GPUs, AMD GPUs, or custom accelerators through pluggable execution providers. ONNX Runtime applies graph optimizations such as constant folding, redundant-node elimination, and operator fusion. Performance relative to a tuned target-specific engine such as TensorRT depends on the workload, execution provider, and hardware.

applies graph optimizations and selects execution providers for the target hardware. This enables single-format deployment across CPUs, GPUs, and specialized accelerators.

Portability still requires inspecting operator placement. An unsupported operation may fall back to another execution provider, introducing host-device transfers and synchronization that preserve correctness while erasing the expected speedup. Profiling must therefore verify provider assignment and operator coverage, not only whether the exported graph runs.

13.9.1.2 Specialized inference engines

When latency or hardware cost binds more tightly than portability, TensorRT²⁹ (NVIDIA GPUs), OpenVINO³⁰ (Intel hardware), and similar engines optimize specifically for their target hardware (NVIDIA 2024c; Intel Corporation 2026b; Chen et al. 2018). They apply aggressive target-specific optimizations that framework-native runtimes may not apply uniformly across hardware.

Layer fusion³¹ combines multiple sequential operations into a single GPU kernel. Consider convolution → batch normalization → rectified linear unit (ReLU) activation. Without fusion, this requires three kernel launches and two intermediate write/read round-trips. Fusion combines all three into one kernel that reads inputs once, computes the combined result in registers, and writes final outputs once.

Kernel auto-tuning selects the fastest algorithm for each operation on the specific GPU. A single convolution can be implemented using dozens of algorithms such as direct, FFT-based, Winograd, and various tiling strategies, each optimal for different input sizes and GPU architectures. Auto-tuning benchmarks each candidate and caches the winner, trading compilation time for runtime performance.

These optimizations can produce substantial speedups over an untuned framework-native baseline, but they require explicit export or compilation and may not support every operation or dynamic shape. The illustrative runtime comparison that follows quantifies one such optimization spectrum.

 Example 13.10: ResNet-50: Runtime comparison

Table 13.19 gives an illustrative batch-one ResNet-50/V100 runtime scenario:

Table 13.19: Illustrative inference runtime comparison: Assumed latency and derived speedup for batch-one ResNet-50 across PyTorch eager, TorchScript, ONNX Runtime, and TensorRT on a V100.

Runtime	Latency	Speedup	Notes
PyTorch (eager)	8.5 ms	1×	Baseline, no optimization
TorchScript	6.2 ms	1.4×	JIT compilation
ONNX Runtime	5.1 ms	1.7×	Cross-platform
TensorRT FP32	2.8 ms	3×	NVIDIA-specific
TensorRT FP16	1.4 ms	6.1×	Tensor Core acceleration
TensorRT INT8	0.9 ms	9.4×	Requires calibration

Systems insight: In this illustrative comparison, TensorRT INT8 provides 9.4× speedup and requires quantization calibration data and NVIDIA-specific deployment.

The optimization-compatibility trade-off is inherent. More aggressive optimization yields better performance yet increases deployment complexity and may introduce numerical differences from training. The choice depends on latency requirements, deployment constraints, and available engineering resources.

After the runtime is chosen, configuration applies the same constraint-first logic. Thread pool sizing controls parallelism for CPU inference: too few threads leave cores idle, while too many cause contention. Memory allocation strategies (preallocated buffers vs. dynamic allocation) trade startup cost against flexibility. Execution provider selection prioritizes which hardware backends handle each operation, and graph optimization level trades compilation time for runtime performance. These settings are not a separate checklist after selection; they are how the selected runtime is made

²⁹ | **TensorRT:** It abandons the portability of general-purpose frameworks by requiring a build phase that optimizes the model for a target GPU architecture (for example, an H100) (NVIDIA 2024c). This hardware lock-in allows aggressive optimizations like layer fusion and precision selection that portable runtimes may not apply uniformly across targets. The resulting nonportable engine can materially reduce latency and therefore the number of GPUs required to meet a throughput target.

³⁰ | **OpenVINO (open visual inference and neural network optimization):** An Intel-oriented inference toolkit that converts, optimizes, and runs models across Intel CPU, GPU, and NPU targets (Intel Corporation 2026b). This direct hardware targeting is an “aggressive” optimization because it abandons some portability that framework-native runtimes must guarantee, allowing it to exploit target-specific kernels and precision choices. The resulting performance gain is workload- and hardware-dependent, but it can make dedicated CPU or edge serving economically viable for smaller and latency-sensitive models.

³¹ | **Layer fusion:** Analogous to loop fusion in compiler optimization, where adjacent loops over the same array are combined to reduce memory traffic. Kernel fusion applies the same principle to GPU operations: sequential kernels that write and re-read intermediate tensors from HBM are merged into a single kernel that can keep data in registers.

honest under production traffic. Production deployments therefore measure configuration impact on latency distributions rather than relying on defaults.

13.9.2 Precision selection for serving

In an illustrative scenario, a team deploying ResNet-50 on V100 GPUs assumes that switching from FP32 to INT8 on the non-Tensor-Core integer path triples throughput and costs less than 0.4 percentage points of accuracy. This example illustrates the direct connection between numerical precision and infrastructure economics. Precision selection connects to the quantization techniques covered in section 10.4. Section D.4 compares the numerical formats (FP32, FP16, BF16, FP8, INT8) and their precision-range trade-offs, and section D.4.2 details the mechanics of symmetric and asymmetric integer quantization. Serving adds runtime concerns such as calibration data availability, layer sensitivity under production inputs, and dynamic precision selection.

13.9.2.1 Precision-throughput relationship

For memory-bandwidth-bound operations, reducing precision can raise throughput by reducing data movement. Equation 13.16 gives the bandwidth-only upper bound from precision reduction:

$$\frac{\text{Throughput}_{\text{INT8}}}{\text{Throughput}_{\text{FP32}}} = \frac{32}{8} = 4 \times (\text{theoretical maximum}) \quad (13.16)$$

Actual speedup depends on the hardware, operator mix, calibration, and kernel alignment, and therefore falls below the bandwidth-only bound whenever another constraint intervenes. Low-precision kernels are most efficient when matrix dimensions fit the target accelerator’s preferred tile shapes. Modern cuBLAS and cuDNN can still use Tensor Cores for many other dimensions, though sometimes less efficiently or with internal padding. Section 11.4.3 provides the detailed Tensor Core architecture that explains these alignment constraints. The precision trade-offs for a standard vision model illustrate how these theoretical limits manifest in practice.

Example 13.11: ResNet-50: Precision trade-offs on V100

Table 13.20 gives illustrative latency, memory, and accuracy assumptions across FP32, FP16, and two INT8 paths for ResNet-50, together with Tensor Core utilization where supported:

Table 13.20: Illustrative precision trade-offs on V100: Assumed latency, memory footprint, and accuracy for ResNet-50 in FP32, FP16, and INT8 (PTQ and QAT), plus Tensor Core utilization where supported.

Precision	Latency	Memory	Accuracy	Tensor Core Util.	Calibration
FP32	2.8 ms	98 MB	76.13%	0%	None
FP16	1.4 ms	49 MB	76.13%	85%	None
INT8 (PTQ)	0.9 ms	25 MB	75.80%	N/A	1,000 samples
INT8 (QAT)	0.9 ms	25 MB	76.05%	N/A	Additional training

Systems insight: In this illustrative scenario, INT8 provides 3.1× speedup and loses 0.33 percentage points of accuracy with PTQ. QAT recovers most of that loss but requires retraining.

13.9.2.2 Precision selection constraints

Precision selection is constrained by layer sensitivity, calibration data, and runtime policy. Not all layers tolerate reduced precision equally. For a scalar uniform b -bit quantizer over the calibrated range $[\alpha, \beta]$, the step and within-range rounding-error bound in equation 13.17 are

$$\Delta = \frac{\beta - \alpha}{2^b - 1}, \quad |\epsilon_{\text{quant}}| \leq \frac{\Delta}{2} \quad (13.17)$$

for values that are not clipped. Values outside the calibrated range are clipped and may incur greater error. Layer-level output or task error also depends on activation distributions and downstream

sensitivity, so it must be measured rather than inferred from a universal norm law. This explains observed patterns where first convolutional and final classification layers are often retained at FP16, while middle layers shown by calibration and task-level evaluation to tolerate reduced precision use INT8.

Post-training quantization adds a data constraint. The calibration dataset determines the scale factors used for INT8 conversion, so it must represent actual serving traffic rather than merely reuse convenient training or validation data. A calibration-serving mismatch can degrade task accuracy, a failure mode revisited in section 13.12.

Advanced serving systems turn precision into a runtime policy. If the system is ahead of its latency SLO, it can use higher precision for better accuracy. For low-confidence INT8 results, it can recompute at FP16. Different customer tiers may receive different precision levels. This pattern enables adaptive quality-latency trade-offs while maximizing throughput during normal operation.

The precision decision has direct infrastructure consequences: under the illustrative $3\times$ throughput assumption, a workload requiring 30 GPUs at FP32 needs 10 at INT8 if all other constraints remain unchanged. The connection between model-level optimization and infrastructure economics is why precision selection cannot be treated as purely a model concern.

13.10 Node-Level Optimization

Runtime selection and precision tuning operate at the model level: they determine *what* computation runs and at *what* numerical format. Between the model and the silicon, however, lies another optimization layer encompassing the mechanics of graph compilation to kernels, byte movement from disk to memory, and CPU-GPU coordination. Node-level work closes the gap between a model benchmark and the complete request path, but its gain depends on which boundary the measured trace identifies as binding.

Consider an image classifier whose model benchmark promises millisecond inference but whose production trace shows a slower request path. Node-level optimization identifies which boundary is wasting time on that machine. The trace usually points to one of four recurring diagnostic boundaries:

- **Graph-to-kernel boundary:** The computation graph has to become a small number of efficient kernels rather than a long sequence of launch overheads.
- **CPU execution boundary:** CPU-side work has to exploit vector units, locality, and runtime libraries rather than scalar Python.
- **Load boundary:** Model bytes have to move from disk into memory fast enough that cold starts do not dominate scale-up events.
- **Host-accelerator boundary:** The host has to keep the accelerator scheduled without gaps caused by preprocessing, transfers, or synchronization.

These are not independent tricks. They are places where a measured trace can explain why a request path is slower than the model benchmark promised.

13.10.1 Runtime graph compilation

Inference engines like TensorRT were introduced in section 13.9. They can outperform an untuned eager baseline because serving often gives the compiler more stable operators and a bounded set of input-shape profiles. When those properties hold, the compiler can spend deployment-time work to remove runtime work; services with highly dynamic control flow or shapes retain fewer of these opportunities.

The first gain is operator fusion, the same kernel-merging optimization section 13.9.1.2 applied to TensorRT. What the static serving graph adds is *when* the fusion happens: because operators and shapes are fixed before any request arrives, the compiler can discover and commit the fused kernels ahead of time rather than rediscovering them at runtime, so no request pays for the analysis.

The same static graph also enables constant folding. If a subexpression depends only on fixed weights or constants, such as $x * (\sqrt{2} / 2)$, the compiler replaces it with the precomputed multiplication $x * 0.707\dots$. This removes work from every request without changing the model's mathematical output.

Memory planning applies the same idea to allocation rather than arithmetic. Since the tensor lifetimes are known, the runtime can precalculate memory offsets and reuse buffers instead of allocating reactively during the request. This eliminates operations while creating a predictable serving path with fewer allocator stalls and less memory fragmentation.

Systems Perspective 13.9: Compilation timing trade-off

Just-in-time compilation waits until the graph is first executed and can specialize to the shapes it observes. That specialization is useful for variable traffic, but it moves compiler work into the serving path and creates a cold-request latency spike.

Ahead-of-time compilation performs the compiler work before deployment. It gives the service a fixed graph and avoids startup compilation latency, at the cost of defining all dynamic shapes explicitly or compiling multiple profiles.

Compilation cost can be paid at two points. JIT pays it in the serving path and risks a first-request latency spike, while ahead-of-time compilation pays it before deployment and requires tighter control over input shapes.

These optimizations lead to a deployment choice. Just-in-time compilation adapts to the shapes observed at runtime, but the first request pays the compilation penalty. Ahead-of-time compilation removes that startup spike by shipping an optimized artifact, but the deployment must explicitly cover every shape profile the service will accept.

13.10.2 CPU inference optimization

CPU inference faces its own optimization landscape, where vectorization, locality, graph optimization, and quantization work together. CPUs remain practical for many inference workloads, especially small models, latency-insensitive jobs, and deployments that cannot keep an accelerator busy. Modern CPUs³² (Intel Xeon, AMD EPYC) contain vector and matrix units such as AVX-512 and AMX, but scalar Python does not exercise them directly. Specialized runtimes such as OpenVINO map neural network operators to optimized kernels that use these instructions (Intel Corporation 2026b).

The next CPU boundary is locality. On multi-socket servers,³³ accessing data on another CPU socket incurs non-uniform memory access (NUMA) latency. An inference server must therefore be NUMA-aware: threads should be pinned to specific cores, and the model weights and input buffers those threads touch should be allocated on the same socket. ML model weights—hundreds of megabytes for a mid-sized network and gigabytes for a large language model—can exceed a CPU’s L3 cache, so the NUMA penalty can be persistent rather than occasional. Inference may repeatedly stream substantial portions of the weight tensor from main RAM, forcing fetches across the slower inter-socket link when locality is not preserved.

For sufficiently small models at batch size one, a CPU can outperform a GPU when accelerator launch, synchronization, and transfer overhead exceed the useful compute time. Model size alone does not determine the crossover; operator mix, resident data, runtime, CPU vectorization, and the latency target must be measured on the candidate hardware.

13.10.3 Model serialization and fast loading

Autoscaling systems are operational control loops that add or remove serving replicas based on load. In those systems, the time to spin up a new node is critical. A major component of “Cold Start” (section 13.6.2) is simply reading the model weights from disk into memory. The choice of serialization format determines how quickly this loading can occur.

PyTorch checkpoints loaded through `torch.load()` use a zip-based serialization container with pickle-encoded metadata; safe loading modes restrict which objects may be reconstructed, but the format still performs framework-specific deserialization. A tensor-oriented format³⁴ can reduce object reconstruction and extra host copies. Memory mapping, introduced in section 13.6.3, further allows tensor bytes to be exposed from the file on demand when the serialized layout is compatible with the CPU tensor view.

³² | **SIMD (single instruction, multiple data):** From Michael Flynn’s 1966 taxonomy of computer architectures, SIMD enables one instruction to operate on multiple data elements simultaneously. Intel’s AVX-512 can process sixteen FP32 values in a 512-bit vector, while AMX adds matrix tile operations. The attainable fraction of peak depends on data type, matrix shape, memory behavior, and implementation, but optimized kernels can be dramatically more efficient than scalar code.

³³ | **NUMA (non-uniform memory access):** Accessing memory local to a CPU socket is faster than accessing memory attached to a different socket; pinning an inference thread to a core is insufficient if its required memory is allocated remotely, forcing traffic across the slower inter-socket link. This failure to co-locate threads and data imposes a workload-dependent latency overhead. The exact penalty depends on socket topology, memory placement, and access pattern, but it can be substantial for memory-bound ML workloads because model weights can exceed L3 cache capacity, causing cross-socket fetches on cache misses.

³⁴ | **Safetensors:** The name emphasizes safety: unlike unrestricted pickle deserialization, the format does not encode executable Python objects (Hugging Face 2026). It stores tensors as contiguous raw bytes with a small header and supports memory-mapped loading. In the local benchmark example, that path is 10× faster than the compared PyTorch checkpoint path; the result is workload- and storage-specific and does not include transfer to a discrete accelerator.

Building on this principle, Safetensors is a tensor format designed for safe, efficient loading. It stores tensor metadata in a small header and tensor values as contiguous bytes, enabling memory-mapped CPU access without reconstructing arbitrary Python objects. Moving those tensors to a discrete accelerator still requires device transfer unless the platform provides a compatible shared-memory path.



Example 13.12: Loading speed: Safetensors vs. Pickle

Scenario: A serverless inference replica loads a 5 GB model checkpoint during cold-start scaling events.

Diagnosis: PyTorch `torch.load()` (Pickle) requires CPU object reconstruction, taking 15 s to initialize. Safetensors memory-mapping bypasses CPU parsing, loading weights in 1.5 s (10× faster).

Systems lesson: Tensor-oriented, memory-mapped formats can reduce CPU reconstruction and eager-copy overhead during cold start. Storage throughput, page faults, runtime initialization, and accelerator transfer still remain in the scale-out budget.

13.10.4 Profiling the serving node

Optimization without measurement is guesswork. The system efficiency metric defined in equation 13.2 provides the target: maximizing the fraction of wall-clock time the accelerator spends on useful computation. Timeline profiling tools like PyTorch Profiler or NVIDIA Nsight Systems (nsys) make that target visible by showing the exact sequence of events on the CPU and GPU.

A useful trace reading is bottleneck-first. Empty spaces in the GPU bar mean idle hardware, usually because the GPU is waiting for CPU preprocessing or disk I/O. Thousands of tiny GPU slivers indicate excessive kernel launches and point toward operator fusion or graph compilation. `MemcpyHtoD` blocks expose host-to-device movement; the diagnostic question is whether those transfers overlap with computation or block it. The timeline therefore converts a vague complaint about slow serving into a concrete boundary in the request path.



Example 13.13: The profiling loop

Scenario: An inference serving node exhibits high p99 tail latency despite low average GPU compute utilization.

Diagnosis: Timeline profiling (Nsight Systems/PyTorch Profiler) reveals large idle gaps between tiny GPU kernels caused by CPU host-side preprocessing bottlenecks and kernel launch overhead.

Systems lesson: Serving node optimization requires empirical timeline profiling. Iteratively identifying trace bottlenecks (kernel launch overhead, host-device copy stalls) guides targeted optimizations like operator fusion and pinned memory.

Table 13.21 is a decision aid rather than a checklist: choose the technique whose target metric matches the measured bottleneck, not the row with the largest displayed gain.

Table 13.21: Node-Level Optimization Impact: Illustrative ranges for matching techniques to measured bottlenecks. Realized gains and implementation cost depend on the model, runtime, hardware, and workload.

Technique	Target Metric	Illustrative Gain	Implement. Cost	Best For
Operator Fusion	Latency & Throughput	2–5×	Medium (Compiler)	Memory-bound layers
INT8 Quantization	Throughput	3–4×	High (Calibration)	Inference-heavy nodes
Graph Compilation	Latency	1.5–3×	Low (One-line)	Static graph models
Zero-Copy Loading	Startup Time	10–50×	Low (File format)	Autoscaling/Cold Start
CPU Pinning	Tail Latency (p99)	20–50% reduction	Low (Config)	Latency-critical apps

This hierarchy of impact guides where to invest engineering effort. A layered checkpoint keeps that prioritization tied to the serving stack, from request transport down to fused kernels.



Checkpoint 13.5: The optimization hierarchy

Optimizing inference follows the request path from the outside in.

The stack has four levels.

- ❑ **System budget:** if a 20 ms SLO spends 6 ms on network and serialization, how much remains for application, model, and kernel work, and what measurements are needed to identify which layer to inspect first?
- ❑ **Application level:** are you batching requests effectively? (Dynamic batching).
- ❑ **Model level:** is the model compiled for the target hardware? (TensorRT, ONNX Runtime).
- ❑ **Kernel level:** are operations fused to minimize memory bandwidth?

13.11 Economics and Planning

Batching, precision tuning, operator fusion, and graph compilation can reduce latency, increase capacity, or lower the cost of a completed inference. Production deployment, however, requires answering a fleet-level question: how many machines, of what type, at what total cost. A team that achieves 1,200 images/second on a V100 still needs to know whether 8 V100s at \$3/hour each or 24 T4s at \$0.53/hour each yields lower cost for its 5,000 QPS target while meeting the same SLO. Serving infrastructure cost grows with sustained request volume (Zhang et al. 2019), while training cost follows its own dataset, model, and optimization budget. The public API price compression shown in figure 13.2 illustrates this pressure: as per-token prices fall, infrastructure efficiency becomes a primary lever for economic viability.

13.11.1 Cost per inference

Cost per inference divides the allocated serving cost over the number of completed inferences that satisfy the service objective. The numerator can include provisioned CPU or accelerator time, memory capacity, data transfer, storage, and orchestration. At high utilization, useful throughput spreads that provisioned cost across many requests. At low utilization, paid capacity sits idle, so the same hourly bill is divided among fewer successful inferences. Applying the framework to ResNet-50 shows how hourly price and sustained valid throughput combine into a unit cost.

13.11.2 GPU vs. CPU economics

In the worked AWS cost analysis in section 13.11.1, GPU instances cost more per hour but deliver much higher parallel throughput. The crossover point depends on model characteristics and latency requirements.

CPU inference makes economic sense for small models with few parameters and simple operations, when latency requirements are relaxed (hundreds of milliseconds acceptable), when request volume is low or highly variable (making GPU reservation wasteful), or when the model's operations do not parallelize well. GPU inference is attractive when models are large with parallel-friendly operations, latency requirements are strict (tens of milliseconds), request volume is high and consistent enough to sustain utilization, and batching can amortize the per-inference overhead of GPU kernel launches.

Beyond steady-state costs, startup time affects scaling economics. Instance launch, image pull, driver initialization, model loading, compilation, and warmup can make accelerator replicas slower to become ready than CPU replicas, although exact times depend on the platform and deployment artifact. If traffic rises faster than useful capacity can come online, latency SLOs can fail despite sufficient eventual capacity.

When accelerator replicas have a long readiness delay, predictive scaling or a warm pool may be more reliable than purely reactive scaling. For a model that runs acceptably on both device types, a hybrid design can use always-on accelerator capacity for baseline load and CPU overflow for spikes, trading a different unit cost and latency profile for faster reserve capacity. This GPU+CPU hybrid is

one instance of the broader patterns cataloged in section 2.10; whether it works depends on model compatibility and identical output semantics across runtimes.



Example 13.14: ResNet-50: Cost analysis

Table 13.22 gives an illustrative AWS-like scenario for hourly cost, throughput, and per-million-image cost when serving ResNet-50:

Table 13.22: Illustrative ResNet-50 cloud inference cost comparison: Assumed hourly cost and saturated throughput, with derived cost per one million images for CPU, T4, and V100 instances.

Instance Type	Cost/Hour	Throughput	Cost per 1M Images
c5.xlarge (CPU)	\$0.17	50 img/s	\$0.94
g4dn.xlarge (T4 GPU)	\$0.53	400 img/s	\$0.37
p3.2xlarge (V100 GPU)	\$3.06	1,200 img/s	\$0.71

Systems insight: The T4 GPU instance has the lowest unit cost under these assumptions. The V100 does not beat the T4 on unit cost at the stated full-utilization rates. Cloud pricing varies by region and changes over time; consult current pricing for production planning.

13.11.3 Capacity planning

The GPU vs. CPU decision establishes the cost per inference, but determining how much infrastructure to provision requires combining cost analysis with the queuing theory foundations from section 13.5. **Capacity planning** translates three inputs into infrastructure specifications: traffic patterns (peak request rate, daily/weekly cycles, growth projections), latency SLOs (p50, p95, p99 targets), and model characteristics (inference time distribution at various batch sizes) (Harchol-Balter 2013).

The worked example in section 13.5 demonstrates the complete workflow: starting from a 50 ms p99 SLO and 5,000 QPS target, deriving the conservative M/M/1 safe utilization threshold of 54 percent from equation 13.6, and determining GPU count with headroom of 62 V100s. This scenario provisions peak load plus 30 percent headroom; production margins depend on the workload and policy. Autoscaling can reduce costs during low-traffic periods while meeting latency objectives during peaks; chapter 14 develops the operational policy layer around these serving calculations. Throughput numbers are meaningful only when coupled with latency guarantees. As the valid-QPS accounting in section 13.5 established, capacity must be sized for requests that actually meet the SLO, not for raw request volume.

13.11.4 Production case study: Serving 8-billion-parameter Llama 3

Beyond resident weights, KV-cache memory is often one of the largest capacity costs in large-language-model serving, and figure 13.7 shows why. Plotted at 70-billion-parameter scale to amplify the effect, the cache grows linearly with context length and batch size until long contexts exhaust the memory left after loading weights and runtime workspace. The 8-billion-parameter Llama 3 profile analyzed in the rest of this section obeys the same physics with more headroom, making it a workload an engineer can fit on one GPU and reason about end to end.

13.11.4.1 Workload profile

The following fixed workload assumptions define the reference case used throughout the latency, memory, and economics calculations in this section; together, they bound the analyses that follow.

- **Model:** 8-billion-parameter Llama 3 (quantized to 4-bit using activation-aware weight quantization (AWQ); see section 10.4 for quantization techniques) (Dubey et al. 2024; Lin et al. 2024).
- **Hardware:** 1 × NVIDIA H100 SXM5 GPU (80 GB HBM3, 3.35 TB/s bandwidth) (Choquette 2023).

- **Request characteristics:** 1,000-token input prompt (Prefill), 256-token generated response (Decode).
- **Target SLOs:** TTFT < 200 ms, TPOT < 20 ms.




Figure 13.7: The KV-Cache Explosion: FP16 KV-cache usage vs. context length for a 70-billion-parameter-class model with GQA. On the modeled 80 GB GPU, 35 GB of INT4 weights leave 45 GB for KV cache before runtime overhead. Batch size 32 reaches that cache budget at roughly 4.3k tokens, forcing a trade-off between batch size and context length.

Longer contexts therefore require smaller concurrent batches or more memory, reducing throughput efficiency.

These assumptions make the case study narrow enough to calculate while preserving the two serving constraints that matter most: the prefill budget for time to first token and the decode budget for each generated token.

13.11.4.2 Latency deconstruction

The end-to-end request latency is governed by the two-phase execution model of autoregressive transformers, applying the TTFT and TPOT metrics defined in section 13.8.1. Prefill determines whether the user sees a prompt response quickly; decode determines whether the generated stream keeps moving after it starts.

🔍 Systems Perspective 13.10: The physics of token generation

Recall the energy-movement invariant quantified in table 4.1: moving a bit is 100–1,000× more expensive than computing on it. In the decode phase, this law determines the physical “cost per word.”

Physics: For batch-one decode with INT4 weights, the weight-dominated arithmetic intensity is approximately 2 FLOPs per parameter divided by 0.5 byte per parameter, or 4 FLOPs/byte. This is far below the H100 ridge point, so the modeled step is memory-bandwidth bound. Larger decode batches reuse the weight stream across sequences and raise arithmetic intensity (Pope et al. 2023). The batch-one bandwidth floor is captured in equation 13.18:

$$T_{\text{token}} \approx \frac{D_{\text{vol}}}{\text{BW}_{\text{memory}}} \quad (13.18)$$

Implication: Every token generation pays an energy cost to move model weights from HBM into compute registers. On an A100 80 GB (2.04 TB/s HBM2e), the same model has a theoretical weight-streaming floor of approximately 1.97 ms per token. When decode remains bandwidth-bound, adding more compute cores yields little latency improvement; faster memory (Physics), smaller models (Algorithm), or better batching and cache management change the bound.

Prefill phase (time to first token) The model processes the 1,000-token prompt in parallel. Under the solver’s stated 40 percent compute-efficiency assumption, the model-side TTFT estimate is 41.9 ms, comfortably within the 200 ms SLO and leaving the remaining budget for network ingress, tokenization, and queueing.

Decode phase (time per output token) The model generates 256 tokens sequentially. At batch one, this phase is memory-bandwidth bound: each step streams the 4 GB weight tensor from VRAM. A batched step amortizes that stream across its active sequences.

To obtain the theoretical floor, we divide the 4 GB weight tensor by the 3.35 TB/s bandwidth: $T_{\text{token}} \approx 1.2$ ms. Adding the solver’s per-layer dispatch tax gives a modeled TPOT of 1.57 ms. Generating all 256 tokens therefore takes $256 \text{ tokens} \times 1.57 \text{ ms} = 0.40$ s.

13.11.4.3 Memory and throughput

With 4-bit weights occupying 4 GB, the remaining ~ 81.9 GB is available for the KV cache and runtime workspace. Each token requires approximately 0.131 MB of FP16 KV cache in the 8-billion-parameter Llama 3 configuration, since grouped-query attention reduces KV-head storage relative to full multi-head attention; AWQ is weight-only and does not make that cache INT4. Dividing the 72 GB reserve by that per-token cost yields capacity for ≈ 0.6 million tokens, so at 1,256 tokens the GPU can hold a concurrent batch size of ~ 469 requests.

13.11.4.4 Unit economics

Consider a representative H100 SXM5 rental cost of approximately \$3/hour. Taking the minimum of isolated prefill capacity and KV-residency capacity gives an optimistic throughput ceiling of 45.2 million tokens/hour and a lower-bound unit cost of \$0.066/million tokens. This calculation does not model prefill/decode resource contention or batched-decode capacity.

This analysis highlights that for LLMs, memory capacity (the size of the KV cache) determines the maximum concurrent residency, while prefill compute and decode bandwidth determine realized token throughput and cost under a specific traffic mix. Memory bandwidth often determines low-batch decode latency; larger batches can shift the limit toward compute or KV-cache traffic.

This case study applies the core principles developed throughout this chapter: latency budgets decompose into prefill and decode phases, queuing theory governs batch sizing and capacity planning, and hardware constraints in the form of memory bandwidth and capacity determine achievable performance and cost. The quantitative framework established here enables principled engineering decisions, but only when applied correctly. Common misconceptions cause even experienced engineers to misapply these principles in practice.

13.12 Fallacies and Pitfalls

Serving inverts training priorities in ways that violate intuitions from batch processing. The nonlinear relationship between utilization and latency, the hidden costs of preprocessing, and the silent failure modes of training-serving skew cause violated SLOs, wasted optimization effort, and accuracy degradation invisible to standard monitoring.

Fallacy: *Reducing model inference latency proportionally reduces user-perceived latency.*

Engineers who optimize model inference expect proportional improvement in user-perceived latency, but serving systems introduce latency sources absent from offline benchmarks. Under load, queuing delay dominates: equation 13.5 shows that at 80 percent utilization with 5 ms service time, average wait time is 20 ms before inference even begins. Reducing inference from 5 ms to 2 ms changes service time but also shifts utilization from 80 percent to 32 percent, reducing queuing wait from 20 ms to 0.9 ms, a $21.2\times$ queuing improvement that dwarfs the 3 ms inference gain. This nonlinear interaction between inference speed and queuing behavior means the *system-level* speedup ($25 \text{ ms} \rightarrow 2.9 \text{ ms}$, or $8.5\times$) far exceeds the *model-level* speedup ($5 \text{ ms} \rightarrow 2 \text{ ms}$, or $2.5\times$). Even a 20 percent service-time reduction can produce a large user-facing improvement at high utilization because it lowers utilization as well as direct compute time. Serving optimization requires analyzing the complete latency budget, including serialization, queuing, preprocessing, and postprocessing, under realistic load conditions rather than profiling inference latency in isolation.

Pitfall: *Running serving infrastructure at high utilization to maximize cost efficiency.*

Teams target 90 percent utilization to minimize idle capacity. In production, latency degrades nonlinearly as utilization approaches capacity. Equation 13.5 shows that at 90 percent utilization, average time in system reaches 10× service time. Moving from 70 percent to 90 percent utilization cuts modeled infrastructure cost by 22.2 percent but triples average latency. For a 5 ms inference service, modeled p99 latency rises from ~76.7 ms to ~230 ms under the M/M/1 assumptions. In this scenario, operating at 60 to 70 percent utilization preserves substantially more latency headroom than 90 percent. A production target must instead be selected from the measured arrival process, service-time distribution, SLO, failure reserve, and cost objective.

Fallacy: *Training accuracy guarantees serving accuracy.*

Engineers assume identical model weights preserve validation-set performance. In production, preprocessing differences can shift inputs outside the training distribution. Section 13.6.1 explains how resize interpolation, numerical precision, and feature timing can create skew despite identical weights. The illustrative scenario here drops from 95 percent to 90 percent, a 5 percentage-point loss; these values are not a reported production incident. Production systems require consistent preprocessing and task-quality monitoring because latency and exception checks cannot detect this failure.

Pitfall: *Using average latency to evaluate serving system performance.*

Engineers monitor average latency because it trends smoothly and is simple to compute. An average, however, cannot show whether the percentile named by an SLO is failing. As section 13.5.5 demonstrates, at 70 percent utilization with 5 ms service time, average latency is 16.7 ms but the modeled p99 reaches 76.7 ms, a 4.6× gap invisible to mean-based monitoring. Services therefore report the relevant latency distribution and SLO percentile rather than the mean alone.

Fallacy: *Larger serving batches always improve throughput without affecting latency SLOs.*

Engineers maximize batch size assuming GPU saturation improves cost efficiency under production load. In serving systems, however, batching introduces a latency-throughput trade-off governed by queuing dynamics absent from offline benchmarks. Accumulating requests into larger batches increases wait time for early arrivals: a batch window of 10 ms means the first request waits 10 ms before inference begins, directly adding to p99 latency. In the representative ResNet-50/V100 scenario in section 13.7.6, increasing batch size from 16 to 32 improves throughput only 12 percent but nearly doubles per-batch inference time from 14 ms to 25 ms, and variable input sizes within a batch can create padding overhead that wastes compute on padding tokens. Section 13.7.3 shows why, for tight p99 targets, larger batch sizes can violate SLOs when batch formation delay plus increased per-batch inference time exceeds the latency budget. Serving batch optimization requires jointly tuning batch size, batch timeout, and concurrency against latency SLOs under realistic traffic patterns, not maximizing throughput in isolation.

Pitfall: *Calibrating quantized models with unrepresentative data.*

Teams may calibrate with a convenient sample even when it does not represent the supported deployment distribution. Post-training quantization determines INT8 scale factors from activation ranges in calibration data, so representative coverage matters. The illustrative scenario here drops from 76.1 percent to 72.9 percent, a 3.2 percentage-point loss; it is not a reported production incident. Effective quantization requires calibration and validation on data representative of the inputs the service is expected to handle.

Fallacy: *Cold start latency only matters for the first request.*

Engineers optimize steady-state latency assuming most requests hit warm instances. Cold starts can compound during traffic spikes, deployments, and recovery. In this illustrative scenario, 10 instances each require 30 s of compilation, totaling 300 instance-seconds; parallel warming makes capacity useful after about 30 s. The assumed cold-request latency is 500 ms versus 5 ms at steady state.

Pitfall: *Scaling without a warm-pool or staged-loading budget.*

Autoscaling policies that count only steady-state replicas underestimate the capacity needed during traffic spikes and deployments. Serving systems need warm pools (pre-initialized spare replicas), staged model loading, or admission control so that new replicas become useful before user requests depend on them. The budget should include compilation, weight loading, cache initialization, and health-check time, because those steps determine whether scale-out adds capacity or adds another source of tail latency.

13.13 Summary

Serving moves a validated model into an operating workload. Online requests add a latency distribution, queuing, traffic variability, and failure reserve to the throughput concerns familiar from training. Little's Law relates stable long-run population, arrival rate, and response time; M/M/1 supplies an intentionally simple model of nonlinear latency growth near saturation. Neither replaces measurement, but together they turn headroom and capacity into explicit hypotheses that a load test can validate.

Effective serving optimization measures the complete request path rather than model inference alone. Protocol and serialization choices matter when transport is a visible fraction of the budget; preprocessing matters when it feeds an accelerator too slowly; batching matters only when arrivals and the SLO permit waiting. Training-serving skew adds a different failure mode: the service can remain fast and error-free while inconsistent transformations silently change its predictions.

The deployment paradigm selected in section 2.1 changes the feasible request path. A cloud API may use dynamic batching under a measured arrival process, synchronized sensors impose a frame deadline, and a single-stream mobile workload may choose batch one to protect responsiveness and energy. The MLPerf scenarios provide standardized approximations of these shapes; production design still begins with the service's own traffic, hardware, and objective.



Key Takeaways: Inverting every training priority

- **Serving is latency economics:** Training rewards throughput over long runs, but serving spends a fixed per-request budget across serialization, preprocessing, queuing, inference, postprocessing, and the network. Optimizing only model latency misses the stages users actually wait on.
- **Utilization turns into waiting:** In the chapter's M/M/1 approximation, average time in system rises from $5\times$ service time at 80 percent utilization to $10\times$ at 90 percent. Headroom keeps modest traffic surges from becoming SLO failures.
- **Fast models reveal pipeline taxes:** Once inference is fast, image decode, tokenization, and other preprocessing can dominate total latency. The binding optimization becomes the request path, not the neural network kernel.
- **Batching follows traffic, not habit:** Poisson-modeled web arrivals can use dynamic batching, synchronized sensors need aligned batches, and single-stream mobile workloads may have little batching opportunity. A useful window converts available latency slack into throughput without violating the SLO.
- **Skew breaks accuracy without errors:** Resize methods, normalization order, calibration data, or feature definitions that differ between training and serving can change the model's effective input distribution. Shared transformations, parity tests, and production-slice monitoring reduce this risk.
- **LLM serving is memory management:** Decode often reads weights from VRAM for every generated token, so token latency is bandwidth-bound unless batching changes the constraint. KV-cache layout, PagedAttention, continuous batching, precision, and runtime choice determine both concurrency and cost per token.
- **Runtime choices become infrastructure bills:** Precision, graph compilation, operator fusion, and serving runtime translate directly into replica count and cost per inference.

Quantization and specialized runtimes can materially reduce required serving capacity when they preserve accuracy and fit the target hardware.


Node-level optimization connects the model's operations to the machine that executes them. Profiling reveals whether compilation, operator fusion, precision, locality, serialization, or loading actually binds the request path. Gains matter only when they survive production input shapes, accuracy checks, and latency variation. Capacity planning converts that valid throughput into replica count and unit cost, as the Llama 3 case study in section 13.11.4 demonstrates. A faster kernel that neither reduces provisioned capacity nor improves the SLO has not produced an economic gain.

Training is judged by useful work completed; an online service is judged by whether requests meet their latency objective. Each stage spends against the same envelope: serialization, preprocessing, the queue, the model, and the response. Data determines input size and arrival shape, the algorithm determines compute and state, and the machine determines service rate and movement costs. Queueing couples all three nonlinearly, so a system within budget at moderate traffic can breach it after a small surge. Serving engineering therefore allocates the budget across the request path, preserves headroom for variability and failure, and verifies that the required fraction of requests finishes before the deadline.

What's Next: From node to factory

A single serving node is fragile: models drift, updates must reach users without interruption, and replicas must behave as one dependable system. Chapter 14 expands the unit of concern from a request to the production factory: CI/CD governs artifacts; model registries and feature stores align serving with training; observability detects latency and accuracy drift; and rollback prevents bad releases from becoming permanent outages.

ML Operations




- 14.1 MLOps Overview
- 14.2 Principles and Foundations
- 14.3 Technical Debt
- 14.4 Development Infrastructure
- 14.5 Production Operations
- 14.6 Design and Maturity Framework
- 14.7 Case Studies
- 14.8 Fallacies and Pitfalls
- 14.9 Summary

Purpose

Why can an ML system be perfectly available and perfectly wrong at the same time?

Conventional monitoring often detects crashes, latency regressions, and availability failures quickly, but semantic errors can remain silent in any software system. ML systems add another silent failure mode. A model experiencing data drift can continue serving ordinary-looking predictions while accuracy degrades, triggering no alerts because every health check (latency, throughput, uptime) remains green. Serving infrastructure gets models into production; operations helps keep them correct once they are there. Model quality can degrade because the world changes. Customer behavior shifts, new product categories appear, seasonal patterns evolve, and the distribution the model learned from diverges from the one it now faces. Drift is not a certainty but a persistent risk for deployed models. Unlike latency or uptime, correctness may become observable only after labels or downstream outcomes arrive, so detecting degradation often requires imperfect signals and explicit investigation rather than a single threshold. Managing it requires continuous monitoring that tracks prediction quality alongside system health, drift signals that trigger investigation, retraining when outcome evidence supports it, and deployment strategies that validate new model versions against production traffic before full rollout. The gap between development and production is not a hurdle cleared once but a condition managed indefinitely. Machine learning operations exists because uptime without measured accuracy may still deliver wrong answers at scale. It makes D·A·M co-design continuous because the data environment can remain a moving target long after initial deployment.





Learning Objectives

- Explain why ML systems can remain available while prediction quality silently degrades under distribution shift
- Diagnose technical debt across data-model, model-infrastructure, and production-monitoring interface boundaries
- Design feature stores, registries, and CI/CD pipelines that preserve training-serving consistency and reproducible rollback
- Apply the retraining staleness model to choose cost-aware retraining triggers and intervals
- Implement layered monitoring for drift, skew, degradation, business metrics, and data freshness
- Compare canary, blue-green, shadow, and rollback strategies for production model release risk
- Evaluate operational maturity and investment using model criticality, operational risk, and organizational readiness

14.1 MLOps Overview

AFTER a model is built, optimized, benchmarked, and served, the system still has to remain correct. A benchmark establishes performance at a point in time; serving infrastructure answers requests in milliseconds. The team deploys to production, and week one looks excellent. The challenge begins in week two.

Data distributions shift, user behavior changes, and the world moves on from the conditions under which the model was trained. ML models that succeed in development can fail to achieve sustained production use when postdeployment ownership and monitoring are absent. The root cause is an operational mismatch: availability-focused monitoring tracks server uptime, request latency, and request success rates, while ML monitoring must also track statistical health, including accuracy over time, input-distribution shift, and per-segment prediction quality. A model can degrade substantially while throwing no exceptions, triggering no infrastructure alarms, and maintaining perfect uptime.

The discipline that makes these invisible failures visible is **machine learning operations (MLOps)**. MLOps synthesizes monitoring, automation, and governance into production architectures that detect degradation, trigger retraining, and maintain system health throughout a model's operational lifetime. It inherits the automation and operations lineage of DevOps (Kreuzberger et al. 2023), but addresses a different failure mode. Conventional services can often be tested against deterministic code paths, while ML systems depend on training data distributions, learned parameters, and environmental conditions that shift continuously.

The week-two problem takes concrete shape in an illustrative deployment scenario. Consider a demand prediction system for a ridesharing service. Initial measurements show 94 percent accuracy, 15 ms p99 latency, and strong performance across test segments. By week four, accuracy has dropped to 88 percent, but the infrastructure metrics show nothing wrong. By week eight, a product manager notices driver dispatch is inefficient; investigation reveals the model has not adapted to a competitor's new promotion that shifted user behavior. The model needed retraining six weeks ago, but no system was watching for this degradation. MLOps provides the framework to detect such drift, trigger retraining, and validate new models before users experience the impact.

The operational mismatch connects directly to the book's analytical foundations. If benchmarking provides the *sensors* for the system, MLOps is the complete *control system*. It closes the verification gap of the verification-gap equation (equation 1.1) by continuously recalibrating against a changing world. MLOps can operationalize a locally fitted degradation equation (equation 1.3) when paired drift and outcome data show that distribution change predicts accuracy loss; divergence alone does not prove inevitable decay. It also formalizes interfaces and responsibilities across traditionally isolated domains (data science, machine learning engineering, and systems operations (Amershi et al. 2019)) through continuous retraining, A/B evaluation, graduated rollout, and standardized artifact tracking that supports reproducibility and auditability.

Deploying, monitoring, and maintaining one production ML system defines the operational unit for this chapter: the **ML node**, a complete system comprising data pipelines, feature computation, model training, serving infrastructure, and monitoring for a single machine learning application. Platform operations at larger scale (managing hundreds of models, cross-model dependencies, multi-region coordination, and organization-wide ML platform engineering) constitute advanced topics that build on these single-model foundations.

The lifecycle of one production ML node starts with the week-two control problem and follows the interfaces that make it observable. Technical debt explains why production ML becomes expensive after the first successful deployment; feature stores, CI/CD pipelines, and experiment tracking then define the infrastructure needed to reproduce data, code, parameters, and configuration. Once those artifacts can be reproduced, monitoring, drift detection, deployment strategy, and incident response keep the model healthy over time. Investment decisions and case studies then show how the same principles look different in an edge wearable and in clinical AI operations.

The single-model operational challenge decomposes into three distinct interfaces. The **data-model interface** is the handoff between data infrastructure and model training; its goal is feature consistency, so training and serving pipelines compute features the same way. The **model-infrastructure interface** is the transition from trained weights to scalable service; its challenge is environment parity, because a model that works in a notebook may fail in production due to version, dependency, or runtime mismatches. The **production-monitoring interface** is the feedback loop that enables self-correction, returning statistical telemetry from production to training because ML systems can degrade through drift without crashing.

Those interfaces determine where the chapter's infrastructure pieces belong. Feature stores stabilize feature computation at the data-model boundary. Model registries and deployment pipelines preserve the model-infrastructure handoff. Drift monitors, retraining triggers, and governance policies close the production-monitoring loop before silent degradation becomes a business failure.

The telemetry¹ flowing through these interfaces provides the data needed for informed operational decisions. That operational scope makes the next task precise: distinguish MLOps from traditional DevOps, identify the foundational principles that govern production decisions, and expose the debt patterns that accumulate when those principles are ignored.

14.2 Principles and Foundations

A production ML release is no longer just a code diff: data distributions, learned parameters, evaluation slices, and monitoring feedback loops all become release objects that can change the system's behavior. MLOps builds on DevOps but addresses these specific demands of ML system development and deployment (Kreuzberger et al. 2023; Amershi et al. 2019). Traditional CI/CD can usually reason about code, configuration, tests, and infrastructure as the primary release objects; ML operations must also manage artifacts whose validity depends on the data and environment that produced them.

DevOps integrates and delivers software systems. MLOps must additionally manage statistical, data-dependent workflows spanning data acquisition, preprocessing, model training, evaluation, deployment, and continuous monitoring through an iterative cycle connecting design, model development, and operations. Trace the infinity-loop structure in figure 14.1 to see how these phases feed back into one another continuously; the loop gives the discipline its operating shape.

Definition 14.1: MLOps

Machine learning operations (MLOps) is the engineering practice of productionizing ML systems through CI/CD, workflow orchestration, reproducibility, versioning, collaboration, continuous training and evaluation, monitoring, and feedback loops (Kreuzberger et al. 2023).

1. **Significance:** The cost of not closing this loop appears as stale predictions, delayed detection, and avoidable recovery work. Drift thresholds, retraining triggers, and mean time to recovery (MTTR) targets are deployment-specific quantities calibrated from the business value of predictions, label delay, validation risk, and retraining cost. The important quantitative habit is not a universal threshold but the control loop: measure

¹ | **Telemetry:** The feedback path that can make model degradation visible before it becomes a business failure; crashes and error codes expose some software failures, but semantic errors can remain silent in any application. ML systems add distribution shift, which can go undetected without statistical telemetry such as feature distributions, prediction confidence, and drift indicators. Availability metrics alone would not expose this degradation.

distribution shift, estimate the cost of staleness, trigger retraining only when expected benefit exceeds validation and rollout risk, and verify the replacement model before promotion.

2. **Distinction:** DevOps can monitor functional correctness and business outcomes as well as *system availability*. MLOps adds statistical controls for *predictive correctness*, which can degrade while uptime, error rate, and latency remain healthy.
3. **Common pitfall:** A frequent misconception is that retraining on new data solves distribution shift. In reality, retraining without first diagnosing which distribution changed (input features ($p(x)$), label relationships ($p(y | x)$), or both) can preserve the underlying failure or amplify sampling and labeling bias. Under covariate shift and support assumptions, reweighting or representative resampling may help; concept drift generally requires current labels and model adaptation.




Figure 14.1: Iterative MLOps Loop: Monitoring and triggering feed production evidence back into design and model development, turning requirements, engineering, testing, and deployment into a continuous operating cycle.

The operational complexity and business risk of deploying machine learning without systematic engineering practices becomes clear in a simple retail scenario. A recommendation model initially improves sales, but silent data drift gradually degrades its predictions. If monitoring tracks system uptime rather than model outcomes, the loss can remain hidden until a later business review. MLOps supplies the controls that make such failures visible before they accumulate into material operational and financial harm.

14.2.1 Foundational principles

That retail deployment illustrates a pattern. Without systematic operational practices, even accurate models fail in production. Read the incident as a debugging sequence. When revenue drops, the first question is whether the team can reconstruct the deployed model, which requires reproducibility. The next question is whether the data pipeline, training job, serving path, and monitoring system have boundaries clear enough to isolate the fault, which requires separation of concerns. If the serving features no longer match the training features, consistency becomes the control that prevents the same incident from returning. If drift begins before users complain, observable degradation is the control that turns silent failure into an alert. Finally, if retraining is possible but expensive, cost-aware automation decides when intervention is worth the operational risk and compute cost. These principles name the controls in the order an operations team needs them.

14.2.1.1 Reproducibility

Reproducibility requires every artifact² that influences model behavior to be versioned and traceable. This principle extends beyond code versioning to encompass data, configurations, and environments. Equation 14.1 expresses this dependency formally:

$$\text{Model Output} = f(\text{Code}_v, \text{Data}_v, \text{Config}_v, \text{Environment}_v; \xi) \quad (14.1)$$

² | **Artifact:** Model weights are realized training outputs; the inputs and randomness that produced them cannot generally be inferred from the parameters. Versioning only code is therefore insufficient, and even versioning code, data, configuration, and environment may not guarantee exact replay when kernels or training are nondeterministic.

where v identifies a version and ξ captures execution randomness. Reproduction requires all four versioned inputs and may also require seeds, data order, deterministic kernels, and compatible hardware. Supporting tools include version control, data versioning, and configuration management.

14.2.1.2 Separation of concerns

Separation of concerns decomposes MLOps systems into distinct functional layers that can evolve independently, as table 14.1 shows:

Table 14.1: MLOps Separation of Concerns: Each layer addresses a distinct responsibility and evolves at different rates across the data, training, serving, and monitoring layers. This separation enables independent scaling and updates, reducing blast radius when changes are required.

Layer	Responsibility	Stability
Data Layer	Feature computation, storage, serving	Changes with data schema evolution
Training Layer	Model development, hyperparameter optimization	Changes with algorithm research
Serving Layer	Inference, scaling, latency management	Changes with traffic patterns
Monitoring Layer	Drift detection, performance tracking	Changes with business requirements

14.2.1.3 Consistency imperative

The separation in table 14.1 enables teams to update serving infrastructure without retraining models, modify monitoring thresholds without redeploying, and evolve data pipelines while maintaining model compatibility. That independence is safe only when training and serving environments process data identically, making training-serving parity a consistency imperative. The financial impact of this inconsistency is captured in equation 14.2:

$$\text{Skew Cost} = \text{Rate}_{\text{skew}} \times Q \times C_{\text{error}} \quad (14.2)$$

where $\text{Rate}_{\text{skew}}$ is the fraction of queries affected by training-serving skew, Q is the total query volume per accounting period, and C_{error} is the average business cost incurred per erroneous prediction.

For a system serving 1,000,000 queries/day with 1 percent skew-induced errors costing \$0.10 each, annual skew cost reaches \$365,000. This quantifies why consistency mechanisms represent investments with measurable returns. These mechanisms include feature stores, shared preprocessing code, and validation checks.

14.2.1.4 Observable degradation

Observable degradation requires ML systems to make silent failures visible through continuous measurement. Model performance can degrade along a continuum rather than fail discretely, and an observed failure's time signature—whether a sudden drop, gradual drift, or subgroup decay—helps determine how it is detected and how the system should respond.

14.2.1.5 Cost-aware automation

Cost-aware automation should balance computational costs against accuracy improvements. Let $\Delta\text{Accuracy}$ be the expected gain in accuracy percentage points, Value per Point the monetary value of one percentage-point gain over the decision horizon, Training Cost the compute and labor cost of a retraining run, and Deployment Risk the expected cost of validation and rollout failure. Equation 14.3 models this trade-off:

$$\text{Retrain if: } \Delta\text{Accuracy} \times \text{Value per Point} > \text{Training Cost} + \text{Deployment Risk} \quad (14.3)$$

The inequality is a decision gate, not an automatic trigger. Retrain only when outcome evidence makes the expected gain exceed both run cost and release risk. Table 14.2 pairs each observable failure signature with an appropriate detector and operational response.

Table 14.2: Degradation Detection Strategies: Four failure signatures map to detectors and operational responses. Sudden drops require immediate diagnosis and may call for rollback, while gradual or subgroup degradation calls for diagnosis before retraining or targeted data collection.

Degradation Type	Detection Mechanism	Response Strategy
Sudden accuracy drop	Threshold alerts	Diagnose; roll back if release-related
Gradual drift	Trend analysis	Diagnose, then retrain if warranted
Subgroup degradation	Cohort monitoring	Targeted data collection
Latency increase	Percentile tracking	Infrastructure scaling

This principle guides the design of retraining triggers, validation thresholds, and deployment strategies examined throughout this chapter. The specific values vary by domain, but the framework for making principled trade-off decisions remains constant. The retraining analysis in section 14.4.2.2 derives the complete economic model with worked examples showing how to calculate optimal retraining intervals. Once the causal chain is clear, the five principles can serve as a compact evaluation framework for tools and practices. The organizing claim of table 14.3 is that each principle is only operational once it is tied to a concrete measurable metric: pairing every principle with its key metric, from artifact hash to net retraining value, is what makes the framework auditable rather than aspirational.

Table 14.3: MLOps Principles Summary: Quick reference for the five foundational principles that guide all MLOps tooling and practice decisions.

Principle	Core Insight	Key Metric
Reproducibility	Version all artifacts	Complete artifact hash
Separation of concerns	Independent layer evolution	Layer coupling score
Consistency	Training equals Serving	Feature skew rate
Observable degradation	Make failures visible	Time to detection
Cost-aware automation	Optimize total cost	Net retraining value

How these principles manifest in practice depends on the workload. A recommendation system drifts daily as user preferences shift; a TinyML model deployed on embedded hardware may run unchanged for months. The monitoring strategy must match the archetype.



Lighthouse 14.1: Monitoring strategy by archetype

The dominant failure modes and monitoring priorities differ across workload archetypes. Table 14.4 compares four representative archetypes by drift pattern, monitoring metric, and example retraining trigger:

Table 14.4: Monitoring Strategy by Workload Archetype: Illustrative starting points for monitoring strategy. Real thresholds must be calibrated to the deployment's label delay, traffic volume, business risk, and alert-fatigue budget.

Archetype	Dominant Drift Pattern	Primary Monitoring Metric	Example Retraining Trigger
ResNet-50 (Compute Beast)	Visual distribution shift (lighting, camera, new object classes)	Accuracy on holdout set (ground truth available)	Accuracy drops > 2% from baseline (~monthly for stable domains)
GPT-2 (Bandwidth Hog)	Vocabulary drift, topic shift, emerging entities	Perplexity on live traffic (no ground truth needed)	Perplexity increases > 10%; new vocabulary detected (~weekly for news domains)
DLRM (Sparse Scatter)	User behavior shift, item catalog churn, cold-start items	CTR/CVR delta vs. historical cohorts	Engagement drops > 5%; catalog refresh (~daily for e-commerce)
DS-CNN (Tiny Constraint)	Acoustic environment change (noise floor shift)	Duty cycle (wakeups/hour) + false positive rate	False wake rate > 1%; battery drain exceeds spec (~quarterly OTA update)

Systems insight: Ground truth availability and physical bottlenecks govern monitoring design. Compute-bound vision models (ResNet-50) bound by $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$ track holdout accuracy; memory-intensive language and recommendation models (GPT-2, DLRM) bound by D_{vol}/BW may rely on perplexity or implicit clicks; microcontroller models (DS-CNN) bound by strict latency L_{lat} and energy limits monitor duty cycles and false wake rates. Retraining cadence is constrained by label delay, deployment access, validation cost, and the rate of observed change.

These principles respond to recurring challenges: concept drift,³ data-quality failures (Schelter et al. 2018), and silent postdeployment degradation. These collectively motivate the specialized tools and workflows distinguishing MLOps from traditional DevOps. The divergence is driven by the silent failure problem introduced at the chapter’s opening: system health cannot be measured by uptime or latency alone. Operational discipline in ML requires monitoring the statistical properties of data distributions and model outputs, shifting the focus from “is the server running?” to “is the system still intelligent?”

Table 14.5 contrasts the objectives, methodologies, primary tools, and typical outcomes of DevOps and MLOps, illustrating how these ML-specific requirements demand distinct operational practices. MLOps coordinates a broader stakeholder ecosystem and introduces specialized practices such as data versioning,⁴ model versioning, and model monitoring that extend beyond traditional DevOps scope. This expanded scope turns model operation into a feedback loop rather than a release pipeline.

Table 14.5: MLOps vs. DevOps: MLOps extends DevOps principles to address the unique requirements of machine learning systems, including data and model versioning, and continuous monitoring for model performance and data drift. MLOps coordinates a broader range of stakeholders and emphasizes reproducibility and scalability beyond traditional software development workflows.

Aspect	DevOps	MLOps
Objective	Streamlining software development and operations processes	Optimizing the lifecycle of machine learning models
Methodology	Continuous Integration and Continuous Delivery (CI/CD) for software development	Similar to CI/CD but focuses on machine learning workflows
Primary Tools	Version control (Git), CI/CD tools (Jenkins, Travis CI), Configuration management (Ansible, Puppet)	Data versioning tools, Model training and deployment tools, CI/CD pipelines tailored for ML
Primary Concerns	Code integration, Testing, Release management, Automation, Infrastructure as code	Data management, Model versioning, Experiment tracking, Model deployment, Scalability of ML workflows
Typical Outcomes	Faster and more reliable software releases, Improved collaboration between development and operations teams	Efficient management and deployment of machine learning models, Enhanced collaboration between data scientists and engineers

 Checkpoint 14.1: The MLOps loop

MLOps is not linear; it is circular.

The feedback cycle

- ☐ **Closing the loop:** how do production metrics (for example, drift alerts) trigger new training cycles?
- ☐ **Retraining trigger:** can you distinguish a drift signal that warrants investigation from outcome evidence that warrants retraining?

The artifacts

- ☐ **Versioning:** are you versioning Data + Code + Model + Environment together?

ML systems add technical-debt sources beyond conventional software. DevOps addresses deployment and scaling, while MLOps also contends with hidden complexity from data dependencies,

³ | **Concept drift:** Concept-drift and data-stream research formalized the problem that a model’s target relationship can change after deployment (Widmer and Kubat 1996; Gama et al. 2014). In adversarial domains such as spam, fraud, and abuse detection, the distribution can actively adapt in response to the model, making continuous monitoring and retraining a structural requirement rather than an operational luxury.

⁴ | **DVC (Data Version Control):** DVC brings Git-like versioning to datasets and model artifacts (Iterative 2024), solving the artifact gap that equation 14.1 formalizes: without data versioning, the Data_v term is unrecoverable, and no combination of code commits can reconstruct the model that was deployed.

model interactions, and evolving requirements. These additional dependencies are sources of technical debt and provide a diagnostic vocabulary for specialized MLOps infrastructure. Understanding boundary erosion reveals why modular pipeline design is necessary. Recognizing correction cascades clarifies why versioning and rollback are essential. Identifying undeclared consumers justifies strict interface contracts. These patterns are the concrete failure modes motivating the infrastructure components that follow.

Each iteration through the loop can introduce data dependencies, model interactions, and configuration drift invisible to standard software testing. Those accumulating costs are technical debt: a framework for converting silent ML failure modes into quantifiable engineering liabilities.

14.3 Technical Debt

The silent failure modes established earlier manifest concretely as technical debt (Sculley et al. 2015): data changes, model interactions, and evolving requirements cause gradual degradation that compounds over time. These failures can accumulate invisibly across multiple system components, demanding engineering approaches that account for statistical behavior and data dependencies. Originally proposed in software engineering in the 1990s,⁵ the technical debt metaphor compares shortcuts in implementation to financial debt, trading short-term velocity for ongoing interest payments in maintenance, refactoring, and systemic risk (Cunningham 1992). In ML, this debt extends beyond code to include “hidden” costs from statistical modeling and data dependencies. Systematic evaluation rubrics, such as the ML Test Score (Breck et al. 2017), provide frameworks for quantifying this debt and assessing production readiness across data, model, and infrastructure components.

⁵ | **Technical debt:** Ward Cunningham’s 1992 WyCash experience report introduced the debt metaphor for expedient code and delayed consolidation (Cunningham 1992); in ML, debt can compound silently through data and model dependencies that code-focused unit tests and reviews may not detect. A pipeline can remain unchanged while the world it models changes. The ML Test Score rubric (Breck et al. 2017) makes this debt explicit through 28 production-readiness tests grouped into data, model, infrastructure, and monitoring sections.




Definition 14.2: Technical debt in ML

Technical debt in machine learning is the accumulating maintenance and change cost created by implicit data dependencies, entangled features, and undeclared consumers in ML systems, where the resulting liabilities may also appear as silent accuracy degradation.

1. **Significance:** Google’s analysis of production ML systems argues that model code is only a small fraction of the surrounding system; the larger operational surface includes data collection, feature extraction, configuration, serving infrastructure, monitoring, and process management (Sculley et al. 2015). ML-specific debt drivers compound this: changing one input feature can silently shift the learned representation of every other feature (entanglement), a model trained to correct another model’s errors creates a fragile dependency chain (correction cascades), and downstream systems consuming model outputs without explicit contracts become undeclared consumers that break silently when the model is updated.
2. **Distinction:** ML technical debt includes system- and data-level dependencies that code review and ordinary code-focused tests may miss. It can slow development and degrade prediction quality even while availability metrics remain healthy.
3. **Common pitfall:** A frequent misconception is that “better code” solves technical debt in ML. In reality, it is a systems architecture problem: the debt accumulates when the assumptions of the training distribution (feature ranges, label meanings, data freshness) are not enforced as runtime contracts at the system boundary.

A break-even calculation makes the cost dynamics of technical debt concrete. Teams often resist automation investment because manual processes seem faster in the short term, but that advantage disappears once repeated manual work exceeds the up-front automation cost. Figure 14.2 places ML code among ten components surrounding a production ML system, making the broader operational surface explicit.

Manual operations hit a capacity ceiling, but the cost problem extends beyond engineering time. ML systems accumulate hidden complexity through specific debt patterns, each emerging from ML’s distinctive reliance on data rather than deterministic logic, statistical rather than exact behavior, and implicit dependencies through data flows rather than explicit interfaces.



Cumulative manual work overtakes the one-time automation investment near week 20.



Figure 14.2: Hidden Infrastructure of ML Systems: ML code is one of ten components arranged around the ML system at the center, alongside data collection, data verification, feature extraction, configuration, machine resource management, serving infrastructure, monitoring, analysis tools, and process management tools. The count is what matters, since modeling occupies one box and the remaining nine are the operational surface a deployment must also carry. Source: (Sculley et al. 2015).



Napkin Math 14.1: The cumulative cost of manual operations

Problem: Why build automated pipelines when manual retraining is faster?

Physics: Manual work accumulates over time.

- Manual retrain: 4 hours of engineering per week.
- Pipeline build: 80 engineering hours (one-time).

Math:

- **Result:** $80 \text{ engineering hours} \div 4 \text{ hours per week} = 20 \text{ weeks}$.
- **Trap:** This assumes the model never changes.
- **Limitation:** Additional features can increase manual effort.
- **Consequence:** After 1 year, manual teams still spend 4 hours per week on maintenance. In this simplified calculation, pipeline teams incur 0 recurring hours.

Context: A central law of systems engineering is that the cost of *maintaining* a system over its lifetime can dominate the cost of *building* it. In ML, technical debt is especially dangerous because it is often *data-driven* rather than *code-driven*: a perfect piece of code can still fail if the data it processes shifts. Measurement is the management boundary: without telemetry, the team cannot tell whether maintenance work is reducing debt or merely hiding it.

Systems insight: Automation addresses capacity limits, not speed alone. A manual team eventually reaches a point where maintaining existing models prevents it from deploying new ones. MLOps responds with systematic observability and automation. Without monitoring infrastructure to make silent failures visible, the team accumulates debt and builds a system that becomes increasingly difficult to manage.

Figure 14.3 maps six patterns across data, model, and infrastructure concerns. The representative examples that follow connect each pattern to the engineering response it demands.

14.3.1 Boundary erosion

The first and often most insidious debt pattern involves the dissolution of system boundaries. In traditional software, modularity and abstraction provide clear boundaries between components,




Figure 14.3: ML Technical Debt Taxonomy: Six debt patterns radiate from hidden technical debt: configuration debt, feedback loops, data debt, pipeline debt, correction cascades, and boundary erosion. The surrounding labels identify their associated failure patterns.

allowing changes to be isolated and behavior to remain predictable. Machine learning systems blur these boundaries for a structural reason: model behavior depends on statistical properties of data flowing through the system rather than on explicit interfaces. A change to upstream data formatting might pass all unit tests while silently degrading downstream model accuracy. This implicit coupling through data, rather than code, creates tightly coupled interactions between data pipelines, feature engineering, model training, and downstream consumption.

This erosion produces **entanglement**: dependencies between components become so intertwined that local modifications require global understanding and coordination. The result is captured by the CACE principle: *Change Anything Changes Everything*. When systems lack strong boundaries, adjusting a feature encoding, model hyperparameter, or data selection criterion can affect downstream behavior in unpredictable ways. For example, changing the binning strategy of a numerical feature may cause a previously tuned model to underperform, triggering retraining and downstream evaluation changes that ripple far beyond the original modification.

The primary defense against boundary erosion is architectural: modularity and encapsulation at the design level. Components with well-defined interfaces allow engineers to isolate faults, reason about changes, and reduce the risk of system-wide regressions. Explicit separation between data ingestion, feature engineering, and modeling logic introduces layers that can be independently validated, monitored, and maintained. Boundary erosion is often invisible in early development because the tight coupling only becomes apparent when a seemingly local change triggers a distant failure. Proactive design decisions that preserve abstraction, systematic testing, and interface documentation provide the most practical defenses against this creeping complexity.

14.3.2 Correction cascades

If boundary erosion describes *how* ML systems lose their structural integrity, correction cascades describe *what* happens when teams attempt repairs. A **correction cascade** occurs when fixing one component introduces problems elsewhere, requiring additional fixes that themselves cause further problems. In ML systems, these cascades are particularly severe because changes propagate through statistical dependencies rather than explicit code paths. Retraining a model to fix one failure mode may degrade performance on previously working cases. Adjusting thresholds to reduce false positives may increase false negatives. Adding features to address edge cases may introduce correlations that destabilize the entire system. Each correction triggers the need for more corrections, creating a cascade that can consume engineering resources far exceeding the original fix.

Figure 14.4 makes the cascade structure visible as a chain of dependent models. Each model is trained to correct the errors of the one before it: model B compensates for model A's residual mistakes, model C compensates for model B's, and so on down the chain. The arrangement holds until an upstream model changes. When model A is retrained to fix one failure mode, every downstream model that was tuned to its previous behavior is invalidated at once, and each must be re-corrected. The red arcs trace this propagation, showing how a repair that looked local reaches across the entire

chain. A change near the top forces the most rework, while one near the bottom is nearly free. For an operations team, one change request can trigger a coordinated retraining of the chain.




Figure 14.4: Correction Cascades: Each model is trained to correct the errors of the model before it, forming a chain of dependencies. Changing an upstream model invalidates every downstream correction at once, turning one local fix into cascading rework across the chain, with each downstream fix demanding its own retraining cycle.

Those arcs matter because they turn local repairs into lifecycle-wide dependencies. Sequential model development is one common source: reusing or fine-tuning existing models accelerates development for new tasks, but it also creates hidden assumptions that are difficult to unwind later. Assumptions embedded in earlier models become implicit constraints for future models, limiting flexibility and increasing the cost of downstream corrections.

Consider a team that fine-tunes a customer churn prediction model for a new product. The original model may embed product-specific behaviors or feature encodings that do not transfer to the new setting. As performance issues emerge, teams may attempt to patch the model, only to discover that the true problem lies several layers upstream in the original feature selection or labeling criteria.

To mitigate correction cascades, teams must balance reuse against redesign. Fine-tuning can reduce computation and data requirements, but it also inherits assumptions from the source model. Training from scratch offers greater control over those assumptions at a much higher data and compute cost. Dataset size alone does not decide between them; task similarity, source-model fit, available data, compute budget, and the cost of inherited dependencies do.

The underlying mechanism is that when model A's outputs influence model B's training data, implicit dependencies emerge through data flows rather than explicit code interfaces. These dependencies are invisible to traditional dependency analysis tools. Preventing cascades requires architectural decisions that preserve system modularity: keeping models loosely coupled, maintaining clear version boundaries, and designing for independent evolution even when reusing components.

14.3.3 Interface and dependency challenges

Boundary erosion and correction cascades both arise when ML systems develop interface dependencies that bypass explicit interfaces. Traditional software dependencies are visible (import statements, API calls, configuration files) and can be analyzed by tools. ML dependencies hide in data. When model A's predictions become features for model B, the dependency exists only in the data pipeline, invisible to code analysis. When a dashboard consumes model outputs to drive business decisions, no interface contract governs the relationship.

Two critical patterns illustrate these challenges. **Undeclared consumers** arise when model outputs serve downstream components without formal tracking or interface contracts. When models evolve, these hidden dependencies break silently. A credit scoring model's outputs might feed an eligibility engine that influences future applicant pools and training data, creating untracked feedback loops that bias model behavior over time. **Data dependency debt** compounds this problem as ML pipelines accumulate unstable and underutilized data dependencies that become difficult to trace or validate. Feature engineering scripts, data joins, and labeling conventions lack the dependency analysis tools available in traditional software development. When data sources change structure or distribution, downstream models fail unexpectedly.

Mitigating these interface challenges requires systematic approaches: strict access controls for model outputs, formal interface contracts with documented schemas, data versioning and lineage tracking systems, and continuous monitoring of prediction usage patterns. The MLOps infrastructure patterns presented in subsequent sections provide concrete implementations of these solutions.

14.3.4 System evolution challenges

The debt patterns in section 14.3.3 describe poor design choices. Even well-designed ML systems face evolution challenges that differ sharply from traditional software.

Feedback loops represent the most subtle evolution challenge: models influence their own future behavior through the data they generate. Recommendation systems exemplify this dynamic: suggested items shape user clicks, which become training data, potentially creating self-reinforcing biases. Operationally, the warning sign is a subgroup error gap that widens across retraining cycles: one cohort receives worse predictions, those predictions reshape future behavior or labels, and the next dataset amplifies the gap. The MLOps lesson is to monitor cohorts before aggregate metrics hide the loop. These loops undermine data independence assumptions and can mask performance degradation for months.

Pipeline and configuration debt accumulates as ML workflows evolve into “pipeline jungles” of ad hoc scripts and fragmented configurations. Without modular interfaces, teams build duplicate pipelines rather than refactor brittle ones, leading to inconsistent processing and growing maintenance burden. Compounding this, rapid prototyping encourages embedding business logic in training code and undocumented configuration changes. While these early-stage shortcuts are necessary for innovation, they become liabilities as systems scale across teams. Managing evolution requires architectural discipline: cohort-based monitoring for loop detection, modular pipeline design with workflow orchestration tools, and treating configuration as a first-class system component with versioning and validation.

14.3.5 Code and architecture debt

Data dependencies and system evolution create debt through implicit coupling. ML systems also accumulate code-level debt patterns that differ from traditional software. Sculley et al. (2015) identify several that deserve explicit attention.

Glue code dominates ML codebases: systems often require substantial integration code to connect general-purpose ML packages to specific data pipelines and serving systems, with the glue constituting up to 95 percent of the codebase while the actual ML code represents only 5 percent. This glue creates tight coupling between package APIs and the surrounding system, meaning that when packages update their interfaces, all glue code must be rewritten. Mitigation requires wrapping ML packages in stable internal APIs and treating external dependencies as substitutable components.

Dead experimental codepaths accumulate as ML development involves extensive experimentation, leaving behind conditional branches for abandoned approaches. Unlike traditional dead code that can be detected statically, experimental ML codepaths often remain “live” because they are controlled by configuration flags rather than compile-time conditions. Over time, these paths increase testing burden and create confusion about which code actually runs in production. Regular code audits with explicit deprecation timelines and feature flag hygiene help manage this debt.

Abstraction debt arises because traditional software engineering relies on well-defined abstractions like functions, classes, and modules, but ML systems lack mature abstractions for key concepts such as the right interface for a “feature” or the right encapsulation for “model behavior.” This absence forces teams to reinvent abstractions or, worse, avoid abstraction entirely. Common patterns such as feature stores (abstracting feature computation), model registries (abstracting model versioning), and prediction services (abstracting inference) reduce per-project abstraction debt when they fit the team’s workflow.

Beyond these patterns, Sculley et al. (2015) identify warning signs, or *common smells*, that indicate accumulating debt: the *Plain-Old-Data Type Smell* (using generic types like strings and floats instead of semantic types that encode meaning and constraints), the *Multiple-Language Smell* (systems spanning Python, SQL, C++, and shell scripts with inconsistent conventions), and the *Prototype Smell* (“temporary” research code that becomes permanent infrastructure without refactoring). Effective

organizations track these smells in code reviews and allocate explicit time for debt reduction, treating technical debt paydown as a first-class engineering activity rather than an afterthought.

14.3.6 Technical debt in practice

The debt patterns described earlier are not theoretical constructs. They have played a critical role in shaping real-world machine learning systems. In practice, unseen dependencies and misaligned assumptions can accumulate quietly, only to become major liabilities over time.

14.3.6.1 Production debt patterns

The first pair exposes coupling through model behavior. YouTube’s recommendation system illustrates how large recommenders learn from noisy user-behavior signals, making ranking objectives, equal per-user weighting, and live A/B evaluation part of the system design rather than offline evaluation details (Covington et al. 2016). Zillow’s home valuation and purchasing workflow exposed the correction-cascade version during its iBuying venture.⁶ Valuation and inventory assumptions propagated into purchasing decisions; later corrections then destabilized inventory and pricing decisions, forcing revalidation and eventually a full rollback when the company shut down the iBuying arm in 2021.

The second pair exposes coupling through ownership and configuration. Safety-critical driving automation illustrates the undeclared-consumer risk from a different direction: when automated-control outputs, driver expectations, and subsystem responsibilities are not specified clearly enough, operational failures can cross component boundaries rather than staying local (National Transportation Safety Board 2017). Facebook’s News Feed iterations show the configuration version of the same governance problem. Rapid experimentation and ranking changes require traceable settings and explicit objectives; otherwise behavioral changes become hard to audit after deployment (Engineering 2016; Mosseri 2018).

These examples are not cautionary tales from careless organizations. They are predictable consequences of deploying probabilistic or automated decision systems without infrastructure that makes coupling visible. YouTube, Zillow, safety-critical driving automation, and Facebook each expose a different debt pattern: feedback loops, correction cascades, undeclared consumers, and configuration sprawl.

Each debt pattern suggests corresponding controls: feature stores can reduce data dependency debt, versioning systems expose configuration debt, CI/CD pipelines constrain pipeline debt, and monitoring systems make feedback loops visible. These are not one-to-one cures but engineering responses to the failure modes diagnosed earlier. Recognizing debt patterns, however, is only half the battle: the organizations in these case studies did not lack talented engineers, but rather lacked systematic controls that could catch specific problems before they compounded. The transition from diagnosis to prevention requires examining each infrastructure component in detail: understanding *what* it does and, more critically, *how* it addresses the failure mode that motivated its creation.

14.4 Development Infrastructure

Development infrastructure turns the debt patterns diagnosed earlier into enforcement points. A feature schema that drifts upstream cannot be repaired by a dashboard alone; it needs a shared contract, a versioned artifact, and a deployment path that rejects incompatible changes before they reach production. Table 14.6 maps each component directly to a foundational principle (section 14.2.1) and a specific failure mode.

Table 14.6: MLOps Infrastructure as Debt Remediation: Each infrastructure component responds to a class of technical debt observed in production ML systems. Feature stores can reduce training-serving skew by reusing feature definitions and values; versioning systems support reproducibility; CI/CD pipelines automate rollout controls; monitoring systems can surface silent degradation before users do.

Infrastructure Component	Principle Implemented	Debt Pattern Addressed
Feature stores	Consistency Imperative	Data dependency debt, training-serving skew
Versioning systems	Reproducibility Through Versioning	Configuration debt, correction cascades
CI/CD pipelines	Cost-Aware Automation	Pipeline debt, boundary erosion
Monitoring systems	Observable Degradation	Feedback loops, silent failures

⁶ | **Zillow iBuying failure:** Zillow reported a plan to wind down Zillow Offers in November 2021, including a Q3 inventory write-down and workforce reductions (Zillow Group 2021). The failure illustrates correction cascade debt at scale: pricing errors, purchasing decisions, and inventory feedback can reinforce one another, creating a loop that no single retraining cycle can break.

Figure 14.5 organizes these components across ML models, frameworks, orchestration, infrastructure, and hardware. Understanding how these layers interact enables practitioners to design systems that systematically address the technical debt patterns identified earlier while maintaining operational sustainability.




Figure 14.5: MLOps Stack Layers: Five columns organize the ML system stack: ML Models/Applications, ML Frameworks/Platforms, Model Orchestration, Infrastructure, and Hardware. MLOps spans orchestration tasks from data management through model serving and infrastructure tasks from job scheduling through monitoring, supporting automation, reproducibility, and scalable deployment.

14.4.1 Data infrastructure and preparation

Reliable machine learning systems depend on structured, scalable, and repeatable data handling. From ingestion to inference, each stage must preserve quality, consistency, and traceability across initial development, continual retraining, auditing, and serving alike. These requirements demand systems that formalize data transformation and versioning throughout the ML lifecycle.

14.4.1.1 Data management

The technical debt patterns described earlier stem largely from poor data management: unversioned datasets create boundary erosion, inconsistent feature computation causes correction cascades, and undocumented data dependencies breed hidden consumers. Data management infrastructure directly addresses these root causes. Building on the data engineering foundations from chapter 4, data collection, preprocessing, and feature transformation become formalized operational processes. Where data engineering focuses on single-pipeline correctness, MLOps data management emphasizes cross-pipeline consistency, ensuring that training and serving compute identical features. Data management thus extends beyond initial preparation to encompass the continuous handling of data artifacts throughout the ML system lifecycle.

Three principles organize the infrastructure that addresses these root causes: consistency, freshness, and quality. Each principle motivates specific tooling rather than the reverse.

The first requirement is data consistency: every artifact influencing model behavior, from raw datasets to engineered features, must be versioned and reproducible. Without versioning, teams cannot trace which data produced which model, making debugging and rollback impossible. The implementation usually combines code versioning, dataset versioning, and durable object storage. DVC (Data Version Control) (Iterative 2024), Git (Torvalds and Hamano 2024), Amazon S3 (Amazon Web Services 2024b), and Google Cloud Storage (Google Cloud 2024b) are examples of that pattern, but the invariant is the important part: raw and processed artifacts must remain addressable by version. Section 14.4.1.3 examines implementation details including Git integration, metadata

tracking, and lineage preservation. At the feature level, a feature store can reduce skew by centralizing feature definitions and serving paths across training and serving pipelines, but parity still requires validation. Uber's Michelangelo platform popularized this pattern inside a large production ML platform, and Feast later made the pattern available as open-source feature-store infrastructure (Hermann and Del Balso 2017; Gojek and Google 2019). Section 14.4.1.2 details implementation patterns for training-serving consistency.

Consistency alone is insufficient if the underlying data is stale. Data freshness ensures that models train and serve on current data rather than outdated snapshots. Automated data pipelines maintain freshness by continuously transforming raw data into analysis-ready formats through structured stages: ingestion, schema validation, deduplication, transformation, and loading. Workflow orchestrators such as Apache Airflow (Apache Software Foundation 2024) and Prefect (Prefect Technologies, Inc. 2024), together with the transformation framework dbt (dbt Labs 2024), make those stages explicit, schedulable, and reviewable as code. Once the pipeline is managed this way, data flows can evolve with model requirements without losing versioning, modularity, or CI/CD integration.

The third pillar, data quality, governs whether the data reaching models is accurate, complete, and consistently labeled. In supervised learning pipelines, labeling quality directly determines model ceilings. Labeling tools such as Label Studio (HumanSignal 2024) support scalable, team-based annotation with integrated audit trails and version histories, capabilities that become essential when labeling conventions evolve over time or require refinement across multiple project iterations.

To illustrate how these three principles reinforce each other in practice, consider a predictive maintenance application in an industrial setting. A continuous stream of sensor data is ingested and joined with historical maintenance logs through a scheduled pipeline managed in Airflow (freshness). The resulting features, including rolling averages and statistical aggregates, are stored in a feature store for both retraining and low-latency inference (consistency). Schema validation, sensor-range checks, missingness tests, and label audits catch malformed or unreliable maintenance records before they reach training (quality), while versioning and model-registry integration preserve traceability from data to deployed model predictions. Data management, organized around these three principles, establishes the operational backbone for model reproducibility, auditability, and sustained deployment at scale.

14.4.1.2 Feature stores

The data dependency debt and training-serving skew patterns described in section 14.3 share a common root cause: inconsistent feature computation across pipeline stages. Consider what typically happens without a feature store: a data scientist computes `user_session_length` in Python for training, while an engineer reimplements the same calculation in Java for serving. Subtle differences emerge: one uses wall-clock time, the other processing time; one includes idle timeouts, the other does not. The model trains on one definition but serves using another, and accuracy degrades silently. Feature stores⁷ address this challenge by providing an abstraction layer between data engineering and machine learning, implementing the consistency imperative through a single source of truth for feature values. In conventional pipelines, feature engineering logic is duplicated or diverges across environments, introducing risks of training-serving skew, data leakage, and model drift.

Feature stores manage both offline (batch) and online (real-time) feature access through a centralized repository. During training, features are computed and stored in a batch environment alongside historical labels. At inference time, corresponding transformation logic is applied to fresh data in an online serving system. This architecture can help models consume consistent features in both contexts, but teams must still test point-in-time correctness, freshness, and online-offline parity. The feature store is, in systems terms, an engineering mechanism that helps control training-serving skew: by centralizing feature definitions and serving them through a shared path, it reduces one source of the pipeline divergence that can cause silent production accuracy loss.

Beyond consistency, feature stores support versioning, metadata management, and feature reuse across teams. A fraud detection model and a credit scoring model may rely on overlapping transaction features that can be centrally maintained, validated, and shared. Integration with data pipelines and model registries enables lineage tracking: when a feature is updated or deprecated, dependent models are identified and retrained accordingly.

⁷ | **Feature store:** Uber's Michelangelo platform described a centralized feature store for sharing and serving features across production models (Hermann and Del Balso 2017). At that scale, the consistency guarantee must hold under an online latency budget: what distinguishes a feature store from a shared library of feature code is that the shared feature path also has to serve fresh features fast enough for real-time inference.

Training-serving skew: Diagnosis and prevention Training-serving skew (defined formally in section 13.6.1) manifests operationally through feature store inconsistencies and pipeline divergence. Table 14.7 summarizes common causes and their detection methods:

Table 14.7: Training-Serving Skew Categories: Each category requires different detection and prevention strategies. Schema and preprocessing skew emerge from code divergence and require parity validation; shared feature logic, including a feature store where appropriate, can reduce them, while data distribution skew requires statistical monitoring against training baselines. Timing skew demands careful analysis of feature freshness between training and serving contexts.

Skew Type	Example	Detection Method
Feature preprocessing	Normalization uses different statistics	Statistical comparison of feature distributions
Missing data handling	Training fills NaN with mean; serving uses 0	Schema validation with explicit null handling
Time-dependent features	Features computed with different time cutoffs	Timestamp validation in feature pipelines
Library version drift	NumPy or Pandas version differences	Environment hash comparison

Training-serving skew case study A practical example illustrates how training-serving skew manifests in production systems. Consider a recommendation system that shows 8 percent accuracy degradation one month after deployment with no model-code changes. Feature distribution comparison reveals that `user_session_length` has a mean of 45 minutes in training but 12 minutes in serving. The root cause is feature-definition skew: the offline training pipeline computes wall-clock duration from the first event to the last event in a session, while the online serving path counts only foreground-active time after idle gaps are removed. As a result, the model learned thresholds tied to a feature definition that production never actually serves.

Feature stores (building on the data pipelines from chapter 4) address this problem by centralizing feature definitions and serving paths for training and serving pipelines. Listing 14.1 demonstrates the pattern: training retrieves point-in-time historical features, serving retrieves current online features, and both calls resolve to the same versioned feature definition rather than duplicated code paths. Parity still requires validation.

Listing 14.1: Feature Store Consistency: Unified retrieval reduces training-serving skew by sharing versioned feature definitions across both pipelines.

```
feature_definitions = registry.load(version="2026-06-01")

training_features = feature_definitions.materialize_historical(
    entities=training_entities,
    at_event_time=True,
    names=["user.session_length", "user.purchase_history"],
)

serving_features = feature_definitions.lookup_online(
    entities=[{"user_id": 12345}],
    names=["user.session_length", "user.purchase_history"],
)

assert training_features.schema == serving_features.schema
assert (
    training_features.definition_hash
    == serving_features.definition_hash
)
```

By defining `session_length` once, training and serving share its logic; parity checks must still verify the resulting values. Centralized feature stores also support feature reuse and metadata tracking, which makes skew easier to detect and correct when a feature definition changes (Hermann and Del Balso 2017; Gojek and Google 2019).

As the consistency imperative quantified (section 14.2.1.3), skew-induced errors at production scale can translate to hundreds of thousands of dollars in annual cost. Feature stores can reduce this recurring leakage through infrastructure investment with measurable returns. The economics become operational only when shared definitions can serve many models and teams.

Example 14.1: Uber Michelangelo feature store

Scenario: Uber’s Michelangelo platform supported production models such as estimated time of arrival and used shared feature infrastructure across offline training and online prediction (Hermann and Del Balso 2017). Separate implementations of the same feature logic would create training-serving skew.

Diagnosis: A shared feature contract reduces the chance that batch and online pipelines interpret a feature differently, while separate offline and online stores still require point-in-time and parity validation.

Systems lesson: Feature stores reduce one major source of training-serving skew by centralizing versioned feature definitions and access paths. They do not eliminate skew: point-in-time correctness, freshness, and online-offline parity still require explicit tests.

Skew detection in CI/CD Automated pipelines should validate feature consistency before deployment. Listing 14.2 shows a function that compares training and serving feature distributions using the Kolmogorov-Smirnov test, rejecting deployment when any feature diverges beyond a threshold. The gate converts an observable distribution mismatch into a CI failure before promotion, but its KS threshold must be calibrated to sample size and feature criticality; a passing statistic does not prove schema, semantic, or point-in-time parity.

Listing 14.2: Feature Skew Validation: This function compares training and serving feature distributions using the Kolmogorov-Smirnov test, rejecting deployment when any feature diverges beyond a configurable threshold.

```
def validate_no_skew(
    training_features, serving_features, threshold=0.1
):
    """Reject deployment if feature distributions diverge."""
    for feature in training_features.columns:
        ks_stat = ks_2samp(
            training_features[feature], serving_features[feature]
        )
        if ks_stat.statistic > threshold:
            raise SkewDetectedError(
                f"{feature}: KS={ks_stat.statistic:.3f}"
            )
```

14.4.1.3 Versioning and lineage

Lineage tracking and versioning implement reproducibility (section 14.2.1), which requires all artifacts influencing model behavior to be versioned. Unlike traditional software, ML models depend on multiple changing artifacts: training data, feature engineering logic, trained model parameters, and configuration settings. MLOps practices enforce tracking of versions across all pipeline components to manage this complexity.

Data versioning allows teams to snapshot datasets at specific points in time and associate them with particular model runs, including both raw data and processed artifacts. Model versioning registers trained models as immutable artifacts alongside metadata such as training parameters, evaluation metrics, and environment specifications. Model registries⁸ provide structured interfaces for promoting, deploying, and rolling back model versions, with some supporting lineage visualization tracing the full dependency graph from raw data to deployed prediction (MLflow Project 2026; Google Cloud 2024d).

These complementary practices form the lineage layer of an ML system. The lineage layer enables introspection, experimentation, and governance by preserving the chain of evidence needed to diagnose a degraded model: whether the input distribution matched training data, whether feature definitions changed, and whether the deployed model version matched the serving infrastructure. By elevating versioning and lineage to first-class citizens in the system design, MLOps enables teams to build and maintain reliable, auditable, and evolvable ML workflows at scale.

⁸ | **Model registry:** Prevents “registry bypass,” the failure mode where the undocumented production model diverges from the trained artifact through different pre-processing, stale serialization formats, or manual hotfixes applied directly to the serving endpoint. Without a registry enforcing versioned, immutable artifacts with queryable metadata and state, rollbacks require locating the correct weights from an ad-hoc artifact store under incident pressure.

14.4.2 Continuous pipelines and automation

Feature stores and versioning systems address data consistency statically: they ensure that features are computed correctly at a point in time. Automation enables these systems to evolve continuously, synchronizing data preprocessing, training, evaluation, and release into integrated workflows that respond to new data, shifting objectives, and operational constraints (Orr et al. 2021; Google Cloud 2026b).

14.4.2.1 CI/CD pipelines

Feature stores and versioning systems address the *data* side of consistency; CI/CD pipelines address the *process* side, ensuring that changes flow through validated stages rather than ad hoc deployments. ML CI/CD pipelines must handle complexity absent from traditional software: data dependencies, model training workflows, and artifact versioning that couple code changes to statistical behavior changes.

A typical ML CI/CD pipeline consists of coordinated stages: checking out updated code, preprocessing input data, training a candidate model, validating performance, packaging the model, and deploying to a serving environment. In some cases, pipelines also include triggers for automatic retraining based on data drift or performance degradation. By codifying these steps, CI/CD pipelines⁹ reduce manual intervention, enforce quality checks, and support continuous improvement of deployed systems.

ML-focused CI/CD layers two tiers of tooling for one reason. A general-purpose CI/CD orchestrator (Jenkins, CircleCI (2024), or GitHub Actions (GitHub, Inc. 2024b)) manages version-control events and execution logic, but the ML layer must additionally version data, gate on model metrics, and trigger retraining. Teams therefore add an ML platform or workflow orchestrator (Kubeflow (Authors 2024), Metaflow (Netflix 2024), or Prefect (Prefect Technologies, Inc. 2024)) that supplies higher-level abstractions for those tasks.

Without this automation, model deployment degrades into a manual, error-prone process: an engineer retrains locally, copies artifacts to a staging server, and promotes to production with no guarantee that the data, code, or hyperparameters match what was validated. The cost of such ad hoc workflows compounds with team size and deployment frequency, producing configuration drift and silent regressions that surface only after the model has served incorrect predictions.

Figure 14.6 shows how a representative continuous-training pipeline moves from dataset ingestion and validation through transformation, training/tuning, evaluation, model validation, and model registration. A retraining trigger initiates the process, while dataset, model, metadata, and artifact repositories preserve the inputs and outputs needed for lineage.




Figure 14.6: ML CI/CD Pipeline: Read the enclosed stages in execution order; the retraining trigger starts a new run, while connected repositories preserve the dataset, model, metadata, and artifacts needed to reproduce and govern it. Adapted from Google Cloud’s MLOps continuous delivery and automation pipeline guidance (Google Cloud 2026b).

To illustrate these concepts in practice, consider an image classification model under active development. When a data scientist commits changes to a GitHub (GitHub, Inc. 2024a) repository,

a Jenkins pipeline is triggered. The pipeline fetches updated data, performs preprocessing, and initiates model training. Experiments are tracked using MLflow (Databricks 2024) which logs metrics and stores model artifacts. After passing automated evaluation tests, the model is containerized and deployed to a staging environment using Kubernetes (Cloud Native Computing Foundation 2024a). If the model meets validation criteria in staging, the pipeline orchestrates controlled deployment strategies such as canary testing (detailed in section 14.4.2.3), gradually routing production traffic to the new model while monitoring key metrics for anomalies. In case of performance regressions, the system can automatically revert to a previous model version.

CI/CD becomes foundational when systems release models repeatedly or carry meaningful production risk, turning ad hoc experimentation into structured, repeatable deployment. Scaling beyond one pipeline requires reusable components and explicit artifact contracts.

Example 14.2: Google TFX production ML pipelines

Scenario: Google's production experience motivated TensorFlow Extended (TFX), a set of reusable components for production ML pipelines (Baylor et al. 2017). As pipelines multiply, ad hoc handoffs make validation, lineage, and reproducibility difficult.

Diagnosis: Each stage needs explicit validation and metadata so a deployed model can be traced to the data, transformations, code, and evaluation that produced it.

Systems lesson: Directed acyclic graph orchestration alone does not establish reproducibility. Standardized components and ML Metadata (MLMD) preserve lineage through recorded inputs and transformations.

14.4.2.2 Training pipelines

CI/CD pipelines orchestrate the overall workflow, but training itself requires specialized infrastructure. Model training, where algorithms are optimized to learn patterns from data, builds on the distributed training concepts covered in chapter 8. Within MLOps, training activities become part of a reproducible, scalable, and automated pipeline supporting continual experimentation and reliable production deployment.

Frameworks such as TensorFlow (Abadi et al. 2016), PyTorch (Paszke et al. 2019), and Keras (Chollet et al. 2024) supply the modular components for building and training models, and the framework-selection principles from chapter 7 carry into production unchanged. This section instead asks which exploratory training logic graduates into a versioned, tested retraining job, and when.

Beyond scalability, reproducibility is a key objective. Training scripts and configurations are version-controlled using tools like Git (Torvalds and Hamano 2024) and hosted on platforms such as GitHub (GitHub, Inc. 2024a). Interactive development environments, including Jupyter (Project Jupyter 2024) notebooks, encapsulate data ingestion, feature engineering, training routines, and evaluation logic in a unified format. In production, notebooks should be treated as exploration harnesses: validated transformations, training code, and evaluation checks must be extracted into versioned, tested modules before they become scheduled retraining jobs.

Notebooks in production CI/CD pipelines assume that code execution is reproducible, but Jupyter notebooks challenge this assumption in subtle ways. While notebooks excel for exploration and prototyping, using them directly in production pipelines introduces operational risks that require mitigation. These considerations are essential for teams transitioning from experimental workflows to production systems.

Reproducibility presents the first challenge. Notebook cells can be executed out of order, creating hidden state dependencies that make results nonreproducible. A common failure mode occurs when a data scientist runs cells one, three, two during development and the resulting model works, but a production pipeline running cells one, two, three fails.

Testing difficulties compound this challenge. Traditional unit testing frameworks do not integrate naturally with notebook structure. Cell-level testing is possible but rarely practiced, leaving notebooks less tested than equivalent Python modules.

Several mitigation strategies address these operational concerns. Papermill enables parameterization and programmatic execution of notebooks, treating them as configurable pipeline stages

(Papermill Project 2026). The nbconvert tool converts validated notebooks to static formats including executable scripts for production execution (Project Jupyter 2026). Cell execution order enforcement tools execute all cells top-to-bottom, rejecting out-of-order dependencies.



Napkin Math 14.2: The cost of silent failures

Problem: Is building an automated drift detection system worth the engineering effort?

Scenario: Consider a product recommendation engine generating \$50M in annual revenue.

Failure mode: A deployment bug causes training-serving skew, dropping recommendation quality by 5 percent. This degrades conversion rate proportionally.

Cost analysis:

1. Manual ops (monthly review):
 - Detection Time: ~4 weeks (28 days).
 - Revenue Loss: $\$50\text{M} \times 0.05 \times 28 \text{ days} / 365 \text{ days} \approx \$191,780.8$.
2. Automated MLOps (daily checks):
 - Detection Time: 1 day.
 - Revenue Loss: $\$50\text{M} \times 0.05 \times 1 \text{ day} / 365 \text{ days} \approx \$6,849.3$.

Systems insight: A single silent failure costs \$184,931.5 more without MLOps. Under this scenario's 4 incidents/year, reducing the time-to-detection (TTD) avoids nearly \$739,726 in annual loss.

The cost calculation makes the notebook-production boundary concrete. Notebooks remain useful for exploration and rapid iteration, but validated logic should move into tested Python modules before it enters production pipelines. The refactoring overhead pays off when it reduces time-to-detection for silent failures; leaving notebook state and preprocessing assumptions implicit turns exploratory convenience into operational risk.

Once training logic is reproducible, automation can standardize the steps around it. MLOps workflows incorporate techniques such as hyperparameter tuning (Ranjit et al. 2019; L. Li et al. 2017), neural architecture search (Elsken et al. 2019), and automatic feature selection (scikit-learn developers 2024a) to explore the design space efficiently. These tasks are orchestrated using CI/CD pipelines, which automate data preprocessing, model training, evaluation, registration, and deployment. For example, a Jenkins pipeline triggers a retraining job when new labeled data becomes available. The resulting model is evaluated against baseline metrics, and if performance thresholds are met, it is deployed automatically.

Public-cloud spending forecasts illustrate the scale of the infrastructure market that supports on-demand training and serving (Gartner 2024). This connects to the workflow orchestration patterns explored in chapter 3, which provide the foundation for managing complex, multi-stage training processes across distributed systems. Cloud providers provision high-performance computing resources on demand, including GPU and Tensor Processing Unit (TPU) accelerators.¹⁰ Depending on the platform, teams construct their own training workflows or rely on fully managed services such as Vertex AI Fine Tuning (Google Cloud 2024c), which support automated adaptation of foundation models to new tasks. Regional hardware availability remains an important consideration when designing cloud-based training systems (Lehdonvirta et al. 2025).

These practices converge on a single operational boundary. Exploratory training logic that proves out in a notebook, once it produces a validated model, is version-controlled and extracted into a scheduled retraining job triggered by data updates or performance monitoring. The exploration harness that made iteration fast is not what runs in production; the tested, versioned module is, and the discipline of that hand-off is what separates a reproducible pipeline from a fragile one. Through standardized workflows, versioned environments, and automated orchestration, MLOps transitions model training from ad hoc experimentation to robust, repeatable systems meeting production standards for reliability, traceability, and performance.

¹⁰ | **Cloud ML training economics:** GPT-3 training cost has been estimated in the millions of dollars when priced in V100 GPU-hours (Li 2020). Fine-tuning costs vary by model size, provider, dataset, and number of training steps. Spot instances and Spot VMs can reduce instance prices but introduce a trade-off: AWS Spot Instances and Google Cloud Spot VMs can be interrupted or preempted, requiring checkpoint-and-resume infrastructure for fault-tolerant training workloads (Amazon Web Services 2026; Google Cloud 2026c).

Retraining decision framework Automated training pipelines introduce a critical decision regarding their execution frequency. Deciding when to retrain a model requires balancing accuracy maintenance against computational costs. Three common strategies exist, each with distinct trade-offs. Table 14.8 provides illustrative schedules across domains, from daily retraining for rapidly shifting ad click prediction to quarterly updates for stable medical imaging applications:

Table 14.8: Illustrative Retraining Schedules by Domain: These represent starting points; actual cadences depend on observed drift rates and business impact, and organizations typically calibrate them through operational experience.

Domain	Illustrative Schedule	Rationale
Ad click prediction	Daily	User interests shift rapidly
Fraud detection	Weekly	Attack patterns evolve continuously
Demand forecasting	Monthly	Seasonal patterns change slowly
Medical imaging	Quarterly	Disease presentations are stable

Those schedules are starting points, not rules. **Scheduled retraining** runs on a fixed cadence, such as daily, weekly, or monthly, regardless of performance metrics. It is simple to implement and guarantees that recent data eventually enters the model, but it can waste compute when the distribution is stable or respond too slowly when a shift happens between calendar runs.

Triggered retraining ties the retraining decision to observed degradation. It optimizes compute cost by retraining only when monitoring detects performance loss or drift beyond thresholds, but it requires robust telemetry and careful calibration to avoid false positives or missed degradation.

Continuous retraining updates the model incrementally as labeled data arrives, either through online learning or periodic micro-updates. This can reduce update lag, but it raises the validation burden because noisy labels or adversarial data can be incorporated before humans have reviewed the shift.

The operating choice therefore depends on four constraints: retraining cost, validation infrastructure, rollback capability, and label availability. Large models may cost tens of thousands of dollars per run; triggered retraining needs outcome evidence or a justified proxy; and every automated update path needs validation and rollback capacity before promotion. Scheduled retraining offers predictable operations, triggered retraining ties investigation to measured change, and continuous retraining reduces update lag at the price of a much tighter validation loop. The appropriate policy depends on the deployment rather than the domain label alone.

The economic and validation constraints become executable in listing 14.3: the decision function schedules a run only when estimated recovered value exceeds retraining, validation, and rollout risk and validation data are fresh.

Listing 14.3: Triggered Retraining Decision: The decision combines quality loss, feature drift, prediction shift, and retraining cost so automatic retraining fires only when the expected benefit exceeds the operational risk.

```
quality_loss = baseline_accuracy - current_accuracy
feature_drift = max(population_stability_index(features))
prediction_shift = distribution_distance(
    baseline_predictions, live_predictions
)

benefit = estimate_value_recovered(
    quality_loss=quality_loss,
    feature_drift=feature_drift,
    prediction_shift=prediction_shift,
)
risk = retraining_cost + validation_cost + rollout_risk

if benefit > risk and validation_data_is_fresh():
    schedule_retraining_run()
```

The gate separates detecting drift from authorizing retraining: drift contributes to estimated benefit, while fresh validation data and positive net value determine whether automation acts. For the complete illustrative scenario evaluated next, the parameters yield an optimal interval near one day.

Quantitative retraining economics The decision to retrain a model balances the cost of accuracy decay against retraining expense. When outcome data support a locally exponential fit, the fitted decay rate gives a measurable timescale.¹¹ Its **half-life**¹² describes when fitted accuracy falls to half its initial value, but it does not determine the economically optimal retraining interval. The derivation in equation 14.4 distinguishes the fitted half-life from the economically optimal retraining interval.

¹¹ | **System entropy:** The fitted decay rate γ can differ substantially across deployments and must be estimated from outcome data. A shorter decay timescale demands more responsive monitoring, but it does not by itself require continuous training; label delay, validation risk, retraining cost, and rollback capacity determine the operational cadence.

¹² | **Half-life** (from nuclear physics, where it measures the time for half of a radioactive sample to decay): In ML operations, the metaphor is a convenient approximation, not a universal law. For the exponential fit in equation 14.5, the half-life is $\ln(2)/\gamma$. Retraining cadence also depends on traffic, value, retraining cost, and risk; seasonal, abrupt, or adversarial change requires a richer drift process.



Napkin Math 14.3: The optimal retraining interval

Problem: How often should the team retrain the model to maximize profit?

Physics: Model accuracy $\text{Accuracy}(t)$ decays at rate γ due to data drift.

- Q : Daily Query Volume (Traffic).
- V : Financial value per query for a unit change in accuracy fraction. With this convention, $V = \$0.50$ means 1 percentage point of accuracy is worth \$0.005 per query.
- C : Fixed cost of a retraining run, including compute and operational overhead.

Formula: The approximation in equation 14.4 gives the optimal retraining interval (T^*) that minimizes the sum of staleness losses and training costs:

$$T^* \approx \sqrt{\frac{2 \cdot C}{Q \cdot V \cdot \text{Accuracy}_0 \cdot \gamma}} \quad (14.4)$$

Math: Consider a lighthouse fraud model ($\text{Accuracy}_0 = 0.95$):

- **Traffic** (Q): 1,000,000 transactions/day.
- **Utility** (V): \$0.50/query for a unit accuracy change.
- **Retraining Cost** (C): \$5,000.
- **Drift Rate** (γ): 2 percent per day.

$$T^* \approx \sqrt{\frac{2 \times 5,000}{1,000,000 \times 0.50 \times 0.95 \times 0.02}} \approx 1 \text{ Day}$$

Systems insight: If traffic is high and accuracy is valuable, the team cannot afford to wait. The pipeline must be automated. If T^* is less than the team's manual deployment time, the system is in a state of permanent technical debt.

The same derivation can be formalized into a framework for calibrating monitoring thresholds based on measurable business impact. This quantitative framework transforms retraining from an ad hoc decision into an engineering optimization, implementing cost-aware automation (section 14.2.1).

The staleness cost function Model accuracy can degrade over time, creating a staleness cost. For economic planning, a team may fit the observable impact over a stable local regime as an exponential decay process. Here γ is a temporal decay rate under the assumption that measured degradation accumulates steadily; it is not determined by distribution divergence alone. The exponential model is a simplification that enables closed-form economic analysis. Let $\text{Accuracy}(t)$ represent accuracy at time t since last training, and Accuracy_0 represent initial accuracy. Equation 14.5 captures this fitted scenario, where the rate γ depends on domain volatility and observed outcomes:

$$\text{Accuracy}(t) = \text{Accuracy}_0 \cdot e^{-\gamma t} \quad (14.5)$$

The cost of staleness accumulates based on query volume Q per time period and the value impact V of a unit change in accuracy fraction. Integrating the instantaneous accuracy loss ($\text{Accuracy}_0 - \text{Accuracy}(t)$) over the retraining interval T yields equation 14.6:

$$\begin{aligned} \text{Staleness Cost}(T) &= \int_0^T Q \cdot V \cdot (\text{Accuracy}_0 - \text{Accuracy}(t)) dt = \\ &Q \cdot V \cdot \text{Accuracy}_0 \cdot \left(T - \frac{1 - e^{-\gamma T}}{\gamma} \right) \end{aligned} \quad (14.6)$$

The integral accumulates cost over time t from 0 to T , and the closed form follows from substituting equation 14.5 for $\text{Accuracy}(t)$.

The retraining cost function Each retraining incurs fixed costs including compute, validation, and deployment overhead. Equation 14.7 decomposes these:

$$\text{Retraining Cost} = C_{\text{compute}} + C_{\text{validation}} + C_{\text{deployment}} + C_{\text{risk}} \quad (14.7)$$

where C_{compute} is the cost of the training run itself, $C_{\text{validation}}$ is the cost of evaluating the new model before release, $C_{\text{deployment}}$ is the cost of rolling it into production, and C_{risk} is the expected cost of potential regression from the new model.

Optimal retraining interval The optimal retraining interval T^* minimizes total cost per unit time, as equation 14.8 shows:

$$T^* = \arg \min_T \frac{\text{Staleness Cost}(T) + \text{Retraining Cost}}{T} \quad (14.8)$$

When we assume $\gamma T \ll 1$, a Taylor expansion of the exponential yields the square-root approximation used in the earlier napkin math calculation. With the parameters in table 14.9, the approximation places the economic optimum near one day; an exact optimum requires minimizing equation 14.8 without the expansion.


Table 14.9: Retraining Decision Parameters: Example values for a fraud detection system processing 1,000,000 transactions daily.

Parameter	Value	Description
Q	1,000,000	Transactions per day
V	\$0.50/query	Value per query for a unit accuracy change
Accuracy_0	0.95	Initial accuracy
γ	0.02	Daily decay rate (2% per day)
Retraining Cost	\$5,000	Total retraining expense

Sensitivity analysis Under the square-root approximation, T^* scales with the square root of these parameters. Table 14.10 shows that a fourfold change in retraining cost, query volume, or decay rate moves the approximated interval twofold.

Table 14.10: Retraining Interval Sensitivity: Under the square-root approximation a fourfold change in an input moves the interval only twofold, and the directions differ. More expensive retraining lengthens the interval, while higher query volume or a faster decay rate shortens it.

Change	Effect on T^*
$4 \times$ retraining cost	$2 \times$ longer interval
$4 \times$ query volume	$2 \times$ shorter interval
$4 \times$ decay rate	$2 \times$ shorter interval



Model limitations This framework provides a first-order approximation that enables principled decision-making, but practitioners should be aware of its assumptions:

- Predictable drift: The exponential decay model assumes drift occurs gradually at a known rate. Sudden data or concept shifts require different detection and response mechanisms.
- Known value function: The model assumes each accuracy point has a quantifiable business value. In practice, this value may be nonlinear or context-dependent.
- Independent retraining cycles: The model treats each retraining decision independently, ignoring potential benefits from continuous learning or transfer across retraining cycles.
- Linear cost scaling: Retraining costs are assumed fixed. In practice, infrastructure costs may vary with compute availability and pricing dynamics.

Despite these limitations, the framework provides a principled starting point for retraining decisions. Parameters improve with calibration against historical data and refinement as operational experience accumulates. By making cost-benefit trade-offs explicit and quantifiable, this framework implements cost-aware automation (section 14.2.1), enabling justified infrastructure investments and monitoring thresholds grounded in measurable business impact.

14.4.2.3 Model validation

Training pipelines produce model candidates; model validation determines which candidates merit production deployment. Unlike research evaluation, where a model that beats a benchmark on a static test set may be considered successful, production validation must assess operational readiness under representative conditions and establish the monitoring and rollback controls needed after the distribution changes.

The evaluation process begins with performance testing against a holdout test set sampled from the same distribution as production data. Core metrics such as accuracy, area under the curve (AUC), precision, recall, and F1 score (Rainio et al. 2024) are computed and tracked longitudinally to detect degradation from data drift (IBM 2024). The three aligned panels in figure 14.7 show this degradation pattern concretely. The top panel presents incoming data samples over time, color-coded by type. The middle panel shows an associated change: a feature distribution (`sales_channel`) gradually shifting from predominantly online to predominantly offline transactions. The bottom panel shows model accuracy declining over the same interval. This visualization captures the core challenge of model validation: the need to monitor *inputs* alongside *outputs* to understand why performance changes.




Figure 14.7: Data Drift Impact: In this illustrative scenario, incoming samples shift from predominantly online to increasingly offline sales while model accuracy declines over the same interval. Monitoring inputs and outcomes together helps determine whether the distribution shift is associated with the loss in accuracy.

Beyond static evaluation, MLOps encourages controlled deployment strategies that simulate production conditions while minimizing risk. One widely adopted method is canary testing (Fowler

2014), in which a new model is deployed to a small fraction of users or queries. During this limited rollout, live performance metrics are monitored to assess system stability and user impact. For instance, an e-commerce platform deploys a new recommendation model to 5 percent of web traffic and observes metrics such as click-through rate, latency, and prediction accuracy. Only after the model demonstrates consistent and reliable performance is it promoted to full production.

Evaluating candidates under identical conditions is the prerequisite for a sound promotion decision, since a candidate that wins only because it was measured against different traffic, features, or time windows tells the team nothing. Cloud ML platforms support this through experiment logging, request replay, and synthetic test-case generation, and tooling such as Weights & Biases (Weights & Biases, Inc. 2024) captures the training artifacts, hyperparameter configurations, and metrics that make those comparisons reproducible and traceable across the training and deployment pipeline.

While automation is central to MLOps evaluation, human oversight remains essential. Automated tests may fail to capture nuanced performance issues such as poor generalization on rare subpopulations or shifts in user behavior. Teams combine quantitative evaluation with qualitative review, particularly for models deployed in high-stakes or regulated environments. This multi-stage evaluation process bridges offline testing and live system monitoring, ensuring models behave predictably under real-world conditions and completing the development infrastructure foundation necessary for production deployment.

14.4.3 Infrastructure integration

The development infrastructure examined earlier addresses two of the three critical interfaces introduced at the chapter's opening. Feature stores and data versioning reduce inconsistency at the data-model interface by centralizing and tracking feature access across training and serving. CI/CD pipelines, model registries, and validation gates address the model-infrastructure interface by automating the transition from trained weights to containerized services with rollback capability.

These represent only two-thirds of the operational challenge, however. A model that passes all validation gates and deploys successfully can still fail silently in production as the world changes around it. The third critical interface, production-monitoring, requires a different set of practices focused not on *building models* but on *keeping them healthy over time*.

14.5 Production Operations

A model that passes every validation gate can still have a half-life. From the moment of deployment, the world may diverge from the training distribution: customers change behavior, competitors launch products, seasons shift, and new edge cases emerge that no test set anticipated. Production operations exist to make this possible decay visible and manageable, implementing the production-monitoring interface through deployment strategies, monitoring, incident response, and governance. The requirements are demanding: handle variable loads, maintain consistent latency, recover gracefully from failures, and adapt to evolving data distributions, all without disrupting service. These practices implement observable degradation at runtime, transforming detected model drift into alerts for investigation before users experience degradation.

14.5.1 Model deployment and serving

Once trained and validated, a model must be integrated into a production environment that delivers predictions at scale. Deployment transforms a static artifact into a live system component, and serving ensures accessibility, reliability, and efficiency in responding to inference requests. Together, these components bridge model development and real-world impact.

14.5.1.1 Model deployment

Consider a fraud detection model that achieves 99.2 percent precision in the development environment. An engineer exports the weights, copies them to a production server, and discovers the model predicts every transaction as legitimate: the production server runs a different version of the feature extraction library, producing inputs the model has never seen. This scenario, frustratingly common, illustrates why deployment is not a file transfer but a systems engineering problem. Packaging, testing, and tracking ML models for reliable production deployment requires treating the model, its

13 | Containerization for ML deployment: Docker (Merkel 2014) packages code with dependencies into portable units; Kubernetes (Burns et al. 2016) orchestrates those units across clusters. Containerization addresses the Environment_o term in equation 14.1 partially by capturing user-space dependencies as a versioned artifact; host kernels, drivers, and hardware remain external dependencies.

14 | Staging validation: ML staging adds probabilistic adequacy checks to conventional functional and operational validation. A model can pass unit tests and still fail in production because the test data does not reflect the deployment distribution, so rollout gates must compare prediction statistics, evaluation slices, and business guardrails against calibrated thresholds rather than rely on unit tests alone.

15 | Shadow deployment: Economically justified when its expected reduction in rollout loss exceeds the cost of shadow infrastructure for duplicated inference without serving results.

16 | Canary deployment: Routes a small fraction of live traffic to a candidate model, using it as a sentinel for production health. The ML-specific challenge is that a “failure” is statistical degradation, not a deterministic crash: detecting a small accuracy difference with high confidence can require thousands of inferences, creating a tension between decision speed and statistical power that determines minimum canary duration.

dependencies, and its configuration as a single deployable unit. One common approach involves containerizing models using container technologies,¹³ ensuring portability across environments.

Production deployment requires frameworks that handle model packaging, versioning, and integration with serving infrastructure. Tools like MLflow and model registries manage these deployment artifacts (A. Chen et al. 2020), while serving-specific frameworks (detailed in chapter 13) handle the runtime optimization and scaling requirements. Before full-scale rollout, teams deploy updated models to staging or QA environments¹⁴ to rigorously test performance.

Shadow deployments¹⁵ duplicate live production traffic asynchronously to a candidate model without returning its inferences to users, enabling latency and output comparison without direct user exposure; accuracy still requires labels or another justified reference. **canary testing**¹⁶ routes a small, incrementally ramped share of user traffic to detect regression under live load. **blue-green deployment**¹⁷ maintains parallel environments for router-level cutover and rapid rollback. Each strategy reduces a different release risk, but none removes the need for tested rollback procedures and production monitoring.

📖 War Story 14.1: The Knight Capital error (2012)

Context: Knight Capital Group was a major market maker in US equities handling over 17 percent of NYSE volume. In August 2012, an operational deployment updated SMARS order-routing software on seven of eight production servers but omitted the eighth (U.S. Securities and Exchange Commission 2013).

Mechanism: The deployment repurposed a binary feature flag that activated legacy Power Peg code on the un-updated eighth server. The stale server entered an un-throttled loop, transmitting parent order requests at ~ 1,400 orders/sec.

Impact: In 45 minutes, the router executed over 4 million unauthorized transactions covering more than 397 million shares, incurring a \$460 million loss. The firm survived only on emergency rescue financing raised days later, and lost its independence in an acquisition within months.

Response: The incident motivates atomic deployment verification across production nodes, explicit feature-flag lifecycle controls, and circuit breakers that limit the damage a faulty release can cause.

Systems lesson: Automated deployment is a physical control problem where configuration drift across heterogeneous nodes causes catastrophic failure. ML deployments share this exact vulnerability: a model registry version mismatch, un-synchronized feature schemas, or partial canary routing can flood production with corrupted inference requests before telemetry alerts fire.

Avoiding the Knight Capital failure mode is exactly why ML deployments stage rollout rather than flip a switch, but staged rollout creates its own problem. When canary deployments reveal problems at partial traffic levels (issues appearing at 30 percent traffic but not at 5 percent), teams need systematic debugging strategies. Effective diagnosis requires correlating multiple signals: performance metrics from chapter 12, data distribution analysis to detect drift, and feature importance shifts that might explain degradation. Teams maintain debug toolkits including A/B test analysis frameworks, feature attribution tools, and data slice analyzers that identify which subpopulations are experiencing degraded performance.

That diagnosis loop must connect directly to the release pipeline. CI/CD integration automates deployment and rollback, but only when rollback is designed as part of the rollout mechanism rather than treated as an emergency script.

Rollback strategies and safety mechanisms Rollback¹⁸ capability is the safety net that enables confident deployment. Without reliable rollback, teams become deployment-averse and slow their iteration velocity. Effective rollback requires planning for three distinct scenarios:

The fastest tier, **immediate rollback**, addresses critical failures detected right after deployment: serving errors, latency spikes, or obvious prediction failures. It requires keeping the previous model version loaded and warm so traffic can switch without cold-start delay. **Rapid rollback** handles

performance degradation detected through canary metrics soon after deployment, which requires model registry integration that keeps previous versions deployable with minimal configuration changes. **Delayed rollback** addresses subtle issues detected through business metrics or user feedback after full deployment, where rollback must account for model-dependent data such as personalization state or cached embeddings accumulated during the new model's operation.

Table 14.11 summarizes implementation patterns for each rollback type:

Table 14.11: Rollback Patterns by Scenario: Each rollback type requires different infrastructure support and state handling strategies. Immediate rollback demands always-warm standbys; delayed rollback may require data migration procedures.

Rollback Type	Trigger	Implementation	State Handling
Immediate	Serving errors, crashes	Hot standby with instant switch	Stateless—no special handling
Rapid	Canary metric degradation	Registry-based redeployment	Clear caches, restart sessions
Delayed	Business metric decline	Full redeployment with migration	Migrate state, replay if needed

Rollback testing Rollback procedures that have never been tested may fail when needed, and the gap is often discovered during an active incident, when cognitive load and time pressure are highest. Failure can stem from configuration dependencies, incompatible state, or unclear ownership. Periodic rollback exercises expose these gaps before they matter. Automated rollback criteria can shorten reaction time, but thresholds must be calibrated to the service and include safeguards against flapping. The restored model must also produce consistent behavior rather than predictions corrupted by stale caches or incompatible feature state. Step-by-step runbooks ensure that the responder need not be the person who designed the deployment.

Stateful vs. stateless rollback ML systems vary in statefulness, affecting rollback complexity:

- **Stateless models:** Classification and regression rollback involves only switching model weights, because each prediction is independent.
- **Stateful models:** Sequential recommendation and conversational systems must consider accumulated user state, and rollback may require session resets, compatibility layers, or state migration. Some systems can capture versioned state checkpoints at deployment boundaries, but clean restoration is an architectural property that must be tested rather than assumed.
- **Models with feedback loops:** Feedback-driven models may not restore previous behavior if training data was contaminated during the problematic deployment window.

A/B testing for model validation A/B testing provides the statistical foundation for deployment decisions by comparing model versions under controlled conditions. Unlike canary deployments (which validate operational stability), A/B tests measure whether a new model improves business outcomes with statistical confidence.

Experiment setup and decision rules A valid A/B test starts with four controls that make the later deployment decision statistically meaningful. The **randomization unit** defines what gets randomly assigned to treatment vs. control. User-level randomization ensures consistent experience but requires larger sample sizes. Request-level randomization enables faster experiments but can confuse users seeing different results.

Under a common-variance Normal approximation, let n be the required users per variant, δ the minimum detectable effect, σ the outcome standard deviation, and $z_{\alpha/2}$ and z_β the positive standard-Normal critical values for the chosen two-sided confidence and power. The worked scenario uses 95 percent confidence and 80 percent power. Equation 14.9 translates those design targets into required traffic before launch.

$$n = \frac{2(z_{\alpha/2} + z_\beta)^2 \sigma^2}{\delta^2} \quad (14.9)$$

For a 2 percent relative lift on a 5 percent baseline conversion rate (5 percent to 5.1 percent) and 80 percent power, we need roughly 745,644 users per variant; with 25,000 users per variant, we could detect only a much larger lift, about 0.5 percentage points absolute.

17 | Blue-green deployment: Maintains two comparable production environments, “blue” (serving current traffic) and “green” (running the candidate), then switches traffic in a routing change once the green environment passes validation. Because rollback is a traffic flip rather than a staged drain, recovery can be faster than a gradual canary for stateless services. The trade-off is extra infrastructure during the switch, so blue-green wins over canary when brief duplicate capacity is acceptable and per-segment statistical validation is expensive.

18 | Rollback (from database transaction management): This “undo” action for deployments is complicated in ML by model-dependent state (for example, cached embeddings), which may be incompatible between model versions. This mismatch can prolong recovery and contribute to deployment aversion and slower iteration.

Guardrail metrics define metrics that must not degrade even if primary metric improves. A recommendation model improving click-through rate by 10 percent while increasing page load time by 500 ms may fail guardrail checks.

For a fixed-horizon test, run until the preregistered sample size and minimum duration are reached, including enough calendar time to capture relevant weekly patterns. Do not stop when significance first appears because repeated peeking inflates false positive rates. Early stopping requires a prespecified sequential design with valid decision boundaries.

Those controls establish the statistical envelope, but ML systems add failure modes that ordinary web experiments can hide. Conversion events may arrive days after prediction, creating delayed feedback: a recommendation shown Monday can drive a purchase Friday, so the attribution window must be part of the test design. Novelty effects can also inflate early performance as users engage with fresh recommendations, which is why mature experiments include a burn-in period before measurement.

Recommendation and ranking systems add interference effects because showing an item to one user can affect what remains available or salient for another, violating the independence assumption behind standard A/B analysis. Segment heterogeneity creates a second analysis problem: an overall neutral result may hide strong positive effects for one cohort and negative effects for another. These complications do not invalidate A/B testing, but they make guardrails, segment analysis, and preregistered decisions part of the experiment rather than after-the-fact interpretation.

Table 14.12 turns those constraints into a deployment decision:

Table 14.12: A/B Test Decision Matrix: Deployment decisions should consider both primary metrics and guardrails. Improvements that come at the cost of guardrail violations require careful trade-off analysis rather than automatic deployment.

Primary Metric	Guardrails	Decision
Significant improvement	All pass	Ship new model
Significant improvement	Some fail	Investigate trade-offs, may need model iteration
No significant change	All pass	Evidence inconclusive; retain current model unless a prespecified equivalence or noninferiority test supports the change
Significant degradation	N/A	Do not ship; investigate root cause

The table is only reliable when the analysis process is disciplined before launch. Teams should preregister expected effects before the test begins and choose the randomization unit, attribution window, guardrails, and minimum runtime before observing outcomes.

The same discipline has to continue during analysis. Sequential testing supports valid interim decisions by predefining when early stopping is statistically allowed, and variance-reduction techniques such as CUPED (Controlled-experiment Using Pre-Experiment Data) reduce metric noise by adjusting outcomes with pre-experiment covariates. Failed experiments should be archived because they encode negative evidence, and the analysis pipeline should be automated so manual spreadsheet work does not become a new source of deployment error.

An A/B decision only matters if the release machinery can promote, hold, or roll back the exact artifact that was tested. Model registries, such as Vertex AI’s model registry (Google Cloud 2024d), act as centralized repositories for storing and managing trained models and versions. Model catalogs serve a different role. Vertex AI Model Garden helps teams discover, test, customize, and deploy Google, partner, and selected open-source models (Google Cloud 2026a). Llama belongs in that model-family and model-catalog context, not in the registry lifecycle claim (Touvron, Martin, et al. 2023).

Inference endpoints carry that tested artifact into live traffic. They typically expose the deployed model via REST APIs for real-time predictions. Depending on performance requirements, teams can configure resources, such as GPU accelerators, to meet latency and throughput targets. Some providers also offer flexible options like serverless¹⁹ or batch inference, eliminating the need for persistent endpoints and enabling cost-efficient, scalable deployments.

To maintain lineage and auditability, teams track model artifacts, including scripts, weights, logs, and metrics, using tools like MLflow²⁰ (Databricks 2024). Together, registries, endpoints, lineage tracking, and distributed orchestration frameworks like Ray²¹ turn A/B outcomes into controlled

19 | Serverless ML inference: The cost-efficiency of this option stems from provisioning compute only upon request and scaling to zero when idle, eliminating the cost of a persistent endpoint. This creates a direct tension with performance targets, as the first request after an idle period incurs a “cold start” latency penalty while the model is loaded into memory. For large models, that delay can be long enough to violate real-time latency budgets unless the service keeps warm capacity or uses a runtime designed for fast loading.

20 | MLflow: An open-source MLOps framework that couples artifact storage (S3/GCS) with metadata DBs (PostgreSQL) to record hyperparameter runs, model binary hashes, and deployment stage transitions. The systems trade-off is centralizing lineage tracking vs. incurring database connection bottlenecks during massive parallel hyperparameter sweeps.

21 | Ray: A distributed computing framework from UC Berkeley (Moritz et al. 2018) that provides a unified task and actor interface, backed by a distributed scheduler and fault-tolerant store. The broader MLOps lesson is that fragmented infrastructure creates translation points where preprocessing logic, normalization constants, tokenizer versions, or artifact formats can diverge silently. Shared execution abstractions can reduce that fragmentation, but training-serving skew still requires explicit consistency checks across data, features, and serving code.

production changes: the tested model can be promoted, observed, and reversed without losing provenance.

14.5.1.2 Model format optimization

A PyTorch model with top benchmark accuracy may serve predictions at 200 ms in production, ten times slower than its service-level objective (SLO). The gap between research frameworks and production serving is often substantial, and format optimization bridges it. Optimized formats can improve latency by converting models into representations tailored for specific hardware, but the gain is workload- and runtime-dependent. The inference runtimes and precision strategies detailed in section 13.9 and section 13.9.2 provide the technical foundations; this section focuses on the operational workflow.

The first operational boundary is representation. Open Neural Network Exchange (ONNX) is a widely used interchange format for model portability, but the choice of runtime and execution provider determines the hardware targets, supported operators, and optimizations available. ONNX Runtime emphasizes a common API across multiple backends, while TensorRT specializes execution for NVIDIA GPUs; actual performance must be measured for the target model and hardware. A typical workflow exports a PyTorch model to ONNX, performs graph cleanup (constant folding and dead-code elimination), fuses compatible operators, selects precision, and validates numerical and task-level behavior against the source model at each step. The representative options in table 14.13 differ in portability, hardware scope, and optimization mechanisms.

Table 14.13: Model Optimization Frameworks: Six representative frameworks differ in supported source formats, target hardware, and optimization mechanisms. TensorRT and TF-TRT are NVIDIA-specific, whereas ONNX Runtime and OpenVINO span multiple processor classes.

Framework	Source Formats	Target Hardware	Key Optimizations
ONNX Runtime	PyTorch, TF, Keras, scikit	CPU, GPU, NPU	Graph optimization, operator fusion, quantization
TensorRT	ONNX, TF, PyTorch	NVIDIA GPU only	Kernel auto-tuning, precision calibration, layer fusion
OpenVINO	ONNX, TensorFlow/TFLite, PyTorch, PaddlePaddle	Intel CPU, GPU, NPU	Model compression, async execution, caching
TF-TRT	TensorFlow	NVIDIA GPU	TensorRT integration within TensorFlow graph
Core ML	TensorFlow, PyTorch	Apple Neural Engine, GPU, CPU	Unified format for Apple devices, on-device inference
TFLite	TensorFlow, Keras	Mobile CPU, GPU, Edge TPU	Quantization, delegate support, model compression

The durable pattern is not the product list but the exchange it exposes: every gain in peak throughput is purchased with some degree of hardware or runtime commitment, so framework choice is a portability versus peak-performance decision before it is a feature comparison. Precision is the second boundary. Quantization reduces model size and increases throughput by using lower-precision arithmetic, but from an operational perspective the key deployment question is not whether INT8 is faster. It is whether the quantized model maintains accuracy under production traffic distributions, not merely calibration datasets. The mechanics of PTQ, QAT, and mixed-precision strategies are covered in chapter 10, with serving-specific precision selection (including dynamic per-request precision) detailed in section 13.9.2.

Production deployment of optimized models therefore requires validation that targets the failure modes optimization can introduce silently. Consider a team that deploys an INT8-quantized model after verifying only throughput improvement: classification accuracy drops on rare but high-value edge cases, and the degradation goes undetected for weeks because aggregate metrics remain within SLO bounds. The first validation layer is numerical equivalence, comparing optimized outputs against the original model on a representative test set with application-specific divergence thresholds. That check is necessary but insufficient because rare inputs, out-of-distribution examples, and subgroup-specific cases can expose quantization artifacts that aggregate test metrics hide.

The second layer is operational validation. Memory footprint must be measured at peak runtime utilization, including dynamic allocations during inference, since some optimizations trade increased

runtime memory for computational speed. Warm-up behavior varies by runtime: Accelerated Linear Algebra may JIT-compile kernels during initial execution, whereas TensorRT selects tactics and builds an engine before deployment; loading that engine can still add startup latency. Runtime version compatibility then closes the loop: deployment configurations need explicit version pinning because even minor runtime changes can affect both performance characteristics and numerical correctness.

14.5.1.3 Inference serving

An optimized model on disk generates no value. It needs infrastructure that accepts requests, runs inference, and returns predictions at scale. The service level agreement (SLA), SLO, and serving architectures detailed in chapter 13 provide the technical foundation; this section focuses on the operational considerations for selecting and managing that infrastructure. At Facebook’s reported scale, serving systems executed tens of trillions of inference operations per day under strict latency targets (Hazelwood et al. 2018), and the gap between a *working* serving system and a *well-operated* one determines whether SLOs are met consistently over months and years.

Production-grade serving frameworks such as TensorFlow Serving (Olston et al. 2017), NVIDIA Triton Inference Server (NVIDIA 2024d), and KServe (KServe Community 2024) provide standardized mechanisms for deploying, versioning, and scaling models. From an operational perspective, the key decision is which framework best fits the deployment context: TensorFlow Serving for TensorFlow-native workflows, Triton for multi-framework GPU serving, and KServe for Kubernetes-native environments requiring scale-to-zero.

Regardless of which serving paradigm is used (online, offline, or near-online, as detailed in section 13.2.2), model inference can be only one part of end-to-end latency. Decomposing the latency budget reveals whether the operational bottleneck lies in the model or elsewhere in the request path.

🔍 Systems Perspective 14.1: The latency budget

A service has a 100 ms p99 SLO, and the model inference budget is 45 ms; the remaining 55 ms must cover every other stage of the request path. Table 14.14 allocates the budget across the request lifecycle.

Table 14.14: Latency Budget Components: Representative allocation of a 100 ms p99 SLO across the request lifecycle. Model inference accounts for 45 percent of total latency, leaving the majority to network, feature retrieval, parsing, postprocessing, and serialization. The optimization-lever column shows where each component can be reduced.

Component	Budget Share	p99 Budget	Optimization Lever
Network RTT	15%	15 ms	Edge deployment, connection pooling
Feature retrieval	25%	25 ms	Feature caching, precomputation
Request parsing	5%	5 ms	Binary protocols (gRPC), schema optimization
Model inference	45%	45 ms	Quantization, batching, model distillation
Postprocessing	5%	5 ms	Async processing, result caching
Response serialization	5%	5 ms	Efficient formats (Protobuf, MessagePack)

Systems insight: Model optimization alone often captures less than 50 percent of the latency opportunity. A model that runs 2× faster reduces this example from 100 ms to 77.5 ms, only 1.3× end-to-end improvement, because inference is 45 percent of total latency.

Systems thinking demands end-to-end analysis. Apply the D·A·M taxonomy to diagnose the root cause across Data (feature extraction overhead, serialization cost), Algorithm (too many layers, unoptimized graph), and Machine (memory bandwidth saturation, thermal throttling). Measure end-to-end performance and optimize the binding bottleneck. If feature retrieval exceeds its budget, no amount of model optimization will achieve the SLO.

Beyond the latency budget, operationalizing serving requires selecting infrastructure techniques for the constraint the budget exposed. Table 14.15 summarizes representative strategies for ML-as-

a-service infrastructure; the organizing question is whether the bottleneck lies in queueing delay, capacity, routing, orchestration overhead, or latency prediction.

Table 14.15: Serving System Techniques: Scalable ML-as-a-service infrastructure relies on techniques like request scheduling and instance selection to optimize resource utilization and reduce latency under high load. For the underlying queuing theory and batching strategies, see chapter 13.

Technique	Description	Example System
Request scheduling & batching	Groups inference requests to improve throughput and reduce overhead	Clipper (Crankshaw et al. 2017)
Instance Selection & Routing	Dynamically assigns requests to model variants based on constraints	INFaaS (Romero et al. 2021)
Predictive Autoscaling	Adds capacity ahead of demand spikes to meet latency SLOs	MArk (Zhang et al. 2019)
Autoscaling	Adjusts model instances to match workload demands	INFaaS
Model Orchestration	Coordinates execution across model components or pipelines	AlpaServe (Li et al. 2023)
Execution Time Prediction	Forecasts latency to optimize request scheduling	Clockwork (Gujarati et al. 2020)

These strategies form the cloud-serving foundation. Edge deployment keeps the same operational goal but changes the constraints: rollback, telemetry, and update control must work on devices with limited power, memory, and connectivity.

14.5.1.4 Edge AI deployment


Consider a smoke detector with an ML model for distinguishing cooking smoke from fire. When this model degrades, an engineer cannot simply SSH into the device, roll back to a previous version, and restart. The device sits on someone’s ceiling with intermittent Wi-Fi, a coin-cell battery, and 256 KB of memory. Every operational assumption from cloud MLOps (instant rollback, centralized logging, real-time monitoring) must be reimagined.

Edge AI represents this shift: machine learning inference occurs at or near the data source rather than in centralized cloud infrastructure (Reddi et al. 2019). Workloads that require low latency, privacy-preserving local processing, intermittent connectivity, or tight energy budgets make edge deployment patterns essential knowledge for MLOps practitioners. The shift introduces three interrelated categories of operational challenges: resource constraints, deployment hierarchy, and update mechanisms.

Resource constraints dominate edge deployment decisions. Edge devices require the aggressive model optimization techniques established in chapter 10 (quantization, pruning, knowledge distillation) to meet the memory and power envelopes of microcontroller-class deployments (Warden and Situnayake 2020; David et al. 2021). Power budgets span four orders of magnitude, from milliwatts for IoT sensors to tens of watts in automotive systems, demanding power-aware inference scheduling and thermal management. Some hard real-time, safety-critical applications impose deterministic timing targets that require worst-case execution time (WCET) analysis under adverse conditions including thermal throttling and memory contention.

These constraints shape a natural deployment hierarchy across three tiers. Sensor-level processing handles immediate data filtering and feature extraction on microcontroller-class devices consuming 1–100 mW. Edge gateway processing performs intermediate inference on application processors with 1–10 W power budgets. Cloud coordination manages model distribution, aggregated learning, and complex reasoning requiring GPU-class resources. This hierarchy enables system-wide optimization: computationally expensive operations migrate upward while latency-critical decisions remain local.

Two deployment contexts deserve specific attention. TinyML targets microcontroller-based inference under tight memory and milliwatt-class power constraints, requiring specialized engines such as TensorFlow Lite Micro and CMSIS-NN (David et al. 2021; Lai et al. 2018). Model architectures must be co-designed with hardware constraints, favoring compact operators, quantization, and pruning strategies whose aggressiveness depends on the device and accuracy target. Mobile AI extends edge deployment to smartphones with moderate compute, using NPUs and GPU compute shaders to meet interactive latency and battery-life constraints through power-aware scheduling.



Edge power budgets span sensors, gateways, and vehicles across orders of magnitude.

Updates and monitoring complete the edge operational picture. Over-the-air (OTA) model updates enable maintenance for physically inaccessible systems. OTA pipelines need secure model distribution, artifact verification, and rollback mechanisms; delta or differential updates can reduce transfer volume when devices need not receive a complete model artifact. Update scheduling must account for device connectivity patterns, power availability, and operational criticality.

Monitoring requires adaptation to resource-constrained environments: lightweight telemetry systems capture essential metrics (inference latency, power consumption, accuracy indicators) while minimizing overhead. Health monitoring tracks device-level conditions (thermal status, battery levels, connectivity quality) to predict maintenance needs. Edge-cloud coordination patterns enable adaptive offloading between tiers based on current load, network conditions, and latency requirements. Feature caching at edge gateways reduces redundant computation, while federated learning lets edge devices contribute model updates rather than raw training records. That protocol choice does not by itself guarantee privacy because updates can still reveal information without additional protections.

Graceful degradation is the defining operational pattern for edge AI. When resources become constrained, systems must maintain essential functionality by reducing model complexity, inference frequency, or feature completeness. This design philosophy must be built in from the start, not bolted on as an afterthought.

Getting models into production is only half the challenge. A successfully deployed model can degrade through drift or data quality issues without triggering any alerts, precisely the silent failure modes that motivated this entire chapter. The monitoring, incident response, and on-call practices that follow close this loop.

14.5.2 Resource management and monitoring

Deployment and serving get models into production. Keeping them healthy requires two complementary disciplines: resource management (provisioning and scaling compute, storage, and networking) and monitoring (observing system behavior and detecting degradation before users notice).

14.5.2.1 Infrastructure management

Three failures illustrate the problem. A model works in staging but fails in production because someone manually provisioned a different GPU type. A training job crashes because a colleague's experiment consumed all available memory. An inference service cannot scale because its resource quotas were set through an informal message six months earlier. These failures share a root cause: infrastructure managed through manual processes rather than code.

Scalable, resilient infrastructure is foundational for operationalizing ML systems, and **infrastructure as code (IaC)** is the practice that makes it reliable. IaC treats infrastructure configuration as software (version-controlled, reviewed, tested, and automatically executed) rather than manually configured through graphical interfaces or command-line tools. This approach brings software engineering discipline to resource management: changes are tracked, configurations can be tested before deployment, and environments can be reliably reproduced.

The specific infrastructure tool matters less than the contract it enforces. Terraform (HashiCorp 2014), AWS CloudFormation (Amazon Web Services 2024e), and Ansible (Hatcher 2024) represent common ways to version infrastructure definitions alongside application code. In MLOps settings, that versioned definition is what lets a team reproduce the GPU type, network policy, storage permissions, and scaling limits used by a training or serving environment across AWS (Amazon Web Services 2024c), Google Cloud Platform (Google Cloud 2024a), Microsoft Azure (Microsoft 2024a), or on-premises infrastructure.

Infrastructure management spans the full ML lifecycle. During training, IaC scripts allocate compute instances with GPU or TPU accelerators, configure distributed storage, and deploy container clusters. Because infrastructure definitions are stored as code, they can be audited, reused, and integrated into CI/CD pipelines ensuring consistency across environments.

Containerization provides the same reproducibility boundary for runtime dependencies. Docker (Merkel 2014) packages the model, libraries, and serving code into an isolated unit, while orchestration systems such as Kubernetes (Cloud Native Computing Foundation 2024a) manage those units

across clusters. The operational value is not the container name; it is the ability to deploy the same artifact repeatedly while resource allocation, scaling, and health management remain explicit.

To handle changes in workload intensity, including spikes during hyperparameter tuning and surges in prediction traffic, teams rely on cloud elasticity and autoscaling.²² Cloud platforms support on-demand provisioning and horizontal scaling of infrastructure resources. Autoscaling mechanisms (Amazon Web Services 2024d) automatically adjust compute capacity based on usage metrics, enabling teams to optimize for both performance and cost-efficiency.

Infrastructure in MLOps is not limited to the cloud. Many deployments span on-premises, cloud, and edge environments, depending on latency, privacy, or regulatory constraints. A robust infrastructure management strategy must accommodate this diversity by offering flexible deployment targets and consistent configuration management across environments.

To illustrate, consider a scenario in which a team uses Terraform to provision a GPU serving node on Google Cloud Platform. The node hosts a containerized TensorFlow model that serves predictions via HTTP APIs, and an autoscaling group adds or removes identical replicas as request load varies. Meanwhile, CI/CD pipelines update the model container based on retraining cycles, and monitoring tools track latency and resource utilization. All infrastructure components, ranging from network configuration to compute quotas, are managed as version-controlled code, ensuring reproducibility and auditability. By adopting Infrastructure as Code, cloud-native orchestration, and automated scaling, MLOps teams can provision and maintain resources required for machine learning at production scale.

Infrastructure as Code addresses how to provision resources; the challenge remains deciding when and how much. ML workloads exhibit qualitatively different resource consumption patterns than stateless web applications: training jobs burst from zero to dozens of GPUs then return to minimal consumption, while inference maintains steady utilization under variable traffic. Training workloads demonstrate bursty requirements that create tension between resource utilization efficiency and time-to-insight. Inference workloads present steadier consumption patterns but with strict latency requirements under variable traffic.

Hardware utilization patterns Provisioning resources is only the first half of the problem; using them efficiently means setting utilization targets that balance cost against reliability, and those targets depend on reading hardware metrics correctly rather than taking them at face value. Understanding hardware utilization patterns is essential for cost-effective ML operations. GPU utilization, like CPU utilization, can mislead operators because throughput also depends on memory, I/O, concurrency, and queuing.

GPU utilization metrics can mislead operators. A high utilization reading might be compute-bound (actively executing tensor operations, the ideal case), memory-bound (waiting for data transfers from GPU memory), or I/O-bound (stalled waiting for input data from CPU or network).

Table 14.16 distinguishes these patterns and their optimization strategies:

Table 14.16: GPU Utilization Patterns: Different utilization signatures require different optimizations. High GPU utilization with low memory bandwidth suggests compute-bound workloads that benefit from parallelism. High memory bandwidth with moderate GPU utilization indicates memory-bound workloads requiring model optimization.

Pattern	GPU Util	Memory bandwidth util.	Optimization Strategy
Compute-bound	>85%	<70%	Larger batch sizes, tensor parallelism within node
Memory-bound	50–85%	>85%	Reduce model size, quantize, optimize memory access
I/O-bound	<50%	<50%	Improve data pipeline, prefetch inputs, use SSDs
Batch-starved	Variable (spiky)	Variable	Dynamic batching, request queuing on single server

Utilization targets by workload Representative utilization targets vary by workload characteristics, reflecting the different latency tolerances and cost sensitivities of each operational mode:

- **Batch training:** Target >80 percent GPU utilization. Lower utilization indicates data pipeline bottlenecks or suboptimal batch sizes. Monitor `gpu_util`, `memory_bandwidth_util`, and `data_load_time`.

²² | **ML autoscaling:** Autoscaling adjusts capacity based on demand signals (Amazon Web Services 2024d), but ML serving adds constraints absent from stateless web services. Autoscaling decisions must account for model loading time (cold-start overhead), GPU memory fragmentation, and batching behavior in addition to CPU utilization. Scaling up too slowly violates latency SLOs; scaling down too aggressively forces repeated cold starts that degrade p99 latency.

- **Online inference:** Target 50–70 percent GPU utilization at p50 load. Reserve headroom (30–50 percent) for traffic spikes. Higher sustained utilization risks latency SLO violations during bursts.
- **Batch inference:** Target >85 percent utilization. Unlike online serving, batch jobs can tolerate queuing delays, enabling maximum hardware efficiency.

Utilization targets are diagnostic starting points, not universal thresholds. The same utilization number can indicate a different bottleneck depending on whether the workload is latency-sensitive serving, throughput-oriented batch inference, or training.

Memory hierarchy effects Model serving performance depends critically on GPU memory hierarchy utilization. Data must flow through multiple memory levels with vastly different bandwidths (section D.3.3 maps the full latency hierarchy across the storage spectrum), as table 14.17 quantifies. L2 is a hardware-managed cache for a small active working set, not a placement tier for selected weights. Full model parameters normally reside in high-bandwidth memory (HBM); larger models may offload parameters or state to host memory or storage, with transfer latency often dominating inference. Section D.1 tabulates the current accelerator specifications and HBM bandwidths these serving numbers draw on, so the capacities and interface bandwidths in table 14.17 trace back to documented per-generation figures, while the on-die L2 cache bandwidth is an approximate value that vendors do not publish directly:

Table 14.17: GPU Memory Hierarchy and Bandwidth: Each level trades capacity for speed. L2 caches a small active working set, while full model parameters normally reside in HBM. Host-memory or storage offload can support models that exceed GPU memory, but transfer latency may dominate inference time.

Memory Level	Bandwidth	Typical Contents
L2 Cache (40 MB on A100)	~3 TB/s	Cached working set
HBM2e GPU Memory (80 GB)	~2 TB/s	Model
PCIe Gen4 x16 to CPU	~32 GB/s	Activations
System RAM (512 GB)	~200 GB/s	Batched inputs
NVMe SSD	~7 GB/s	Model swap

For large language model (LLM) serving on a single GPU or server, the KV-cache (storing attention keys and values for each token) often becomes the memory bottleneck; vLLM’s PagedAttention design was motivated by this serving pressure (Kwon et al. 2023). For a Llama 2 70-billion-parameter-style grouped-query attention model (Touvron, Martin, et al. 2023) with 80 layers, 8 KV heads, a 4,096-token context, and FP16 cache entries, each active sequence stores about 1.3 GB of KV cache. Eight concurrent sequences therefore consume about 10.7 GB before scheduler headroom, fragmentation, or activations, limiting how many requests a single node can batch together. Monitoring KV-cache utilization on each serving node enables capacity planning. Near the memory limit, the scheduler may queue or reject requests, reduce batch size, or shorten the admissible context, each with a different latency or service-quality cost.

Cost-per-inference tracking Let Hourly GPU cost be the billed cost of the serving GPU for one hour and Inferences per hour its sustained throughput over the same interval. Equation 14.10 converts those hardware metrics into business-relevant cost-per-inference:

$$\text{Cost per 1K inferences} = \frac{\text{Hourly GPU cost} \times 1000}{\text{Inferences per hour}} \quad (14.10)$$

For an illustrative GPU at \$3/hour processing 50,000 inferences/hour, cost is \$0.06/1K inferences. Track this metric over time; changes may reflect pricing, workload mix, or efficiency.

14.5.2.2 Model and infrastructure monitoring

Infrastructure management provisions resources; monitoring observes their behavior. Unit tests can verify deterministic components but cannot by themselves establish population-level predictive performance, which must be estimated statistically. Monitoring implements observable degradation

(section 14.2.1), transforming this theoretical limitation into operational practice. Once monitoring surfaces a symptom—a latency SLA miss, throughput below target, or memory creep—section A.5.1 maps that symptom to its dominant D·A·M term and tells the operator which optimizations will move the binding constraint and which will be wasted on serving infrastructure. Without continuous monitoring and the deeper **observability** it enables (the ability to infer internal state from outputs), a deployed model is a black box slowly drifting toward irrelevance.

Effective monitoring spans both model behavior and infrastructure performance. On the model side, teams track metrics such as accuracy, precision, recall, and the confusion matrix (scikit-learn developers 2024b) using live or sampled predictions to detect whether performance remains stable or begins to drift. A critical constraint is the **drift detection delay**, which determines how quickly statistical monitoring can confirm that degradation has occurred. The speed of detection depends on traffic volume. A short sample-rate calculation makes that constraint visible.

The sample-rate calculation for the drift detection delay exposes a fundamental asymmetry: statistical tests require enough labeled samples to achieve power, and low-traffic systems may wait days or weeks before accumulating sufficient evidence. This latency gap is not an engineering shortcut that better tooling can close; it is a consequence of finite sample rates colliding with the statistical power requirements of hypothesis testing. The practical implication is that monitoring systems must distinguish between drift that alters the input distribution (detectable without labels) and drift that changes the decision boundary itself (detectable only after ground truth arrives).

Napkin Math 14.4: The drift detection delay

Problem: A 95 percent-accurate model may fall by 5 percentage-point. The monitoring policy budgets 1,000 labeled examples; the exact requirement depends on the test, power, label noise, and dependence. How long does collection take?

Math:

1. Evidence budget: 1,000 labeled examples.
2. High-rate case: At 1 labeled outcome per second, collection takes 1,000 seconds $\approx$ 16.7 minutes.
3. Low-rate case: At 100 labeled outcomes per day, collection takes 10 days.

Systems insight: Detection is bounded by labeled-outcome arrival, which may lag far behind traffic. Low-volume or delayed-label domains may need days or weeks, so high-stakes systems supplement statistical monitoring with proactive model audits.

Production ML systems require rigorous **lineage tracking** within a model registry. Where practical, each registered model version should link immutable identifiers or digests for the code, training data, configuration, and weight artifact that produced it. This chain supports audit and rollback, but it does not guarantee exact reproducibility when nondeterminism, external services, drivers, or hardware remain outside the recorded environment. Production systems face two dimensions of model drift²³ that monitoring must distinguish, although they can occur together. Concept drift occurs when the relationship between features and targets changes, so $p(y | x)$ shifts. Data drift²⁴ refers to a change in the input distribution $p(x)$. In applications such as self-driving cars, this may result from seasonal changes in weather, lighting, or road conditions.

Both forms of drift motivate a formal definition:

Definition 14.3: Data drift

Data drift is a change in the input distribution $p(x)$. Covariate shift is the special case in which $p(x)$ changes while $p(y | x)$ remains stable. The broader drift taxonomy from section 4.5.3 also includes concept drift, in which $p(y | x)$ changes; both dimensions can occur together.

1. **Significance:** It violates the identical-distribution assumption and can cause accuracy to erode as the distributional divergence ($\mathcal{D}(P_t \| P_0)$) grows. When paired drift and outcome



Drift detection speed is bounded by the sample rate.

23 | **Drift detection delay:**

A change in $p(x)$ can be estimated from sufficient input samples without labels, although detection is not immediate. A change in $p(y | x)$ generally requires labeled outcomes, which in high-stakes domains (medical diagnosis, fraud detection, legal decisions) can take days, weeks, or months. Proxy metrics such as prediction-confidence distributions and output entropy can provide imperfect early warnings that trade false-alarm rate for detection speed.

24 | **Covariate shift:**

Importance-weighting corrections for covariate shift assume that the support of the training distribution covers the deployment distribution: every deployment input could have appeared in training, just with different probability. When deployment contains genuinely out-of-distribution inputs (new product categories, new demographics, adversarial inputs), the correction can fail and the model can produce confidently wrong outputs with no warning signal, making support coverage the hidden assumption that determines whether drift correction or full retraining may be required.

measurements support it locally, teams may fit $\text{Accuracy}(t) \approx \text{Accuracy}_0 - \lambda \cdot \mathcal{D}(P_t \| P_0)$ with λ fit per deployment; this is a heuristic, not a general law. Under covariate shift and support assumptions, importance weighting or retraining with current labeled data may recover performance.

2. **Distinction:** Model decay describes observed quality decline over time; the model code need not change. Data drift is one possible external cause, alongside concept drift and changes in the surrounding pipeline.
3. **Common pitfall:** Monitoring model outputs alone may miss input drift or confuse benign output changes with quality loss. Input feature statistics ($\mathcal{D}(P_t \| P_0)$ via PSI, KS tests, or Wasserstein distance/Earth Mover's Distance (EMD)) can provide earlier warning than labeled performance because ground-truth feedback is often delayed.

Because of drift, a deployed model may behave like *decaying inventory* rather than *static software*. **Statistical drift risk** is modeled locally by $(\text{Accuracy}(t) \approx \text{Accuracy}_0 - \lambda \cdot \mathcal{D}(P_t \| P_0))$, with λ fitted from paired divergence and outcomes; this is not a universal monotonic law. Monitoring must track both before degradation compounds into business impact.

The Rotting Asset Curve plot (figure 14.8) puts this entropy into perspective by contrasting two maintenance strategies. The orange sawtooth pattern represents scheduled retraining: accuracy resets at a fixed interval, whether the model is still healthy or has already fallen below the drift threshold. This approach is simple but can be both wasteful and late because the calendar, not observed degradation, drives the update. The green line represents an outcome-triggered response: monitored quality crossing the threshold triggers diagnosis, and retraining follows when labeled evidence and the diagnosed cause support it. The decay rate and intervals are illustrative.




Figure 14.8: The Rotting Asset Curve: An illustrative exponential accuracy decay under two maintenance policies, both drawn as sawtooths. The scheduled policy resets every 90 days regardless of accuracy and dips below the drift threshold before each reset, while the triggered policy resets only on crossing that threshold and so holds a higher floor. Neither cadence is universal.

The two curves turn drift from an abstract statistical problem into an operations policy choice. Scheduled retraining is easy to plan but can retrain too early or too late; trigger-based retraining requires stronger telemetry but aligns intervention with observed degradation.

14.5.3 Layered monitoring and drift quantification

The statistical drift risk established earlier shows that distribution divergence can accompany accuracy decay, but labeled outcomes are needed to determine the direction and magnitude. Quantifying that relationship requires two layers of telemetry: infrastructure metrics that reveal whether the serving system itself is the bottleneck, and distribution and outcome metrics that connect data movement to model performance. Gradual long-term degradation is particularly insidious because it can evade coarse detection thresholds: small day-to-day changes in a quality metric can com-

pound into material degradation over a year without tripping monthly alerts. Seasonal patterns compound this complexity. A model trained in summer may perform well through autumn but fail in winter conditions it never observed. Detecting such gradual degradation requires multi-timescale monitoring: performance baselines across multiple time horizons (daily, weekly, quarterly), sliding window comparisons that detect slow trends, and seasonal performance profiles that account for cyclical patterns.

Systems Perspective 14.2: Iron law in production monitoring

These utilization patterns map directly to the iron law of ML systems (section 1.7). Monitoring reveals which term dominates:

- **Compute-bound** (high GPU utilization, low memory bandwidth utilization): Limited by $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$. Optimize kernels, use Tensor Cores, or upgrade hardware.
- **Memory-bound** (moderate GPU utilization, high memory bandwidth utilization): Limited by D_{vol}/BW . Optimize with quantization, pruning, or batching.
- **I/O-bound** (low GPU utilization, low memory bandwidth utilization): Limited by data pipeline latency. Fix the DataLoader, not the model.

The iron law doubles as a diagnostic framework for production systems. When latency SLOs are violated, the monitoring dashboard indicates which term to investigate.

The first layer is infrastructure-level monitoring, which tracks indicators such as CPU and GPU utilization, memory and disk consumption, network latency, and service availability. GPU utilization alone is incomplete; correlate it with memory-bandwidth and pipeline metrics to distinguish compute-, memory-, and I/O-bound behavior. Power-efficiency metrics (for example, inferences per joule or FLOP/s/W, depending on workload) add a cost-normalized view that enables mixed-workload scheduling for both economic and environmental impact.

Napkin Math 14.5: The economics of observability

Trade-off: “Measure everything” is physically impossible at scale. Let N_{series} be the expanded series count, f_{sample} the sampling frequency per series, B_{sample} the bytes per sample, and $T_{\text{retention}}$ the retention duration; $C_{\text{ingest/sample}}$ and $C_{\text{storage/byte-time}}$ are provider-specific ingestion and storage prices. Equation 14.11 combines these terms into an operating cost rate:

$$\text{Cost rate} \approx N_{\text{series}} f_{\text{sample}} (C_{\text{ingest/sample}} + B_{\text{sample}} T_{\text{retention}} C_{\text{storage/byte-time}}) \quad (14.11)$$

The product is an operating cost rate, not a one-time setup cost.

Both data volumes follow from the same two-step arithmetic. At full fidelity, every request’s telemetry enters the stream: $1\text{M req/s} \times 1\text{ KB per request} = 1\text{ GB/s}$. Retaining one in 60 request traces reduces average ingest to 16.7 MB/s ; batching sampled traces every 60 s changes delivery timing, not the average byte rate. Table 14.18 places the two regimes side by side with their cost impact.

Table 14.18: Trace Sampling vs. Observability Cost: Retaining one in 60 request traces yields a 60× data-volume difference at 1M req/s , assuming 1 KB per request . Batch intervals do not cause the reduction.

Sampling	Granularity	Data Volume (1M req/s)	Cost Impact
All traces	Micro-bursts	~1 GB/s	High (Requires dedicated cluster)
1-in-60 traces	Trends	~16.7 MB/s	Lower (Policy-dependent)

Systems insight: Retain 1 percent of successful requests and 100 percent of errors. Monitor aggregate counters every 1 s and high-cardinality sketches every 60 s; section 14.5.3.1 shows how to budget the infrastructure.

25 | **Prometheus:** Prometheus periodically scrapes targets and stores time series for aggregation. Scrape and rule-evaluation intervals jointly affect alert latency and may miss short excursions. Monitoring per-accelerator thermals allows precise workload routing at higher data cost, while server-level aggregates can mask component-level throttling.

Thermal monitoring integrates into operational scheduling decisions, particularly for sustained high-utilization deployments where thermal throttling can degrade performance unpredictably. An MLOps monitoring dashboard can incorporate thermal headroom metrics to guide workload distribution across available hardware and reduce thermal-induced performance degradation that can violate inference latency SLAs. Tools such as Prometheus²⁵ (Cloud Native Computing Foundation 2024b), Grafana (Labs 2024), and Elastic (Elastic NV 2024) are widely used to collect, aggregate, and visualize these operational metrics. These tools often integrate into dashboards that offer real-time and historical views of system behavior.

Collecting all of these signals at production scale introduces its own cost constraints. Those constraints force engineering teams to make deliberate trade-offs between monitoring granularity and infrastructure expense.

The remaining question is how alerting mechanisms convert statistical signals into actionable responses before the cost of silent failure exceeds the cost of intervention. Proactive alerting mechanisms notify teams when anomalies or threshold violations occur. A sustained drop in model accuracy may trigger drift investigation; infrastructure alerts can signal memory saturation or degraded network performance. The design of these alerts determines the gap between when degradation begins and when an engineer acts on it, and that gap translates directly into business impact at scale. Alerts should preserve the model version, affected segment, and triggering window so responders can reproduce the failure state.

Example 14.3: Recommendation monitoring at scale

Scenario: Consider a high-throughput streaming recommendation service whose latency remained within its SLO while recommendation quality declined as user behavior changed.

Diagnosis: Traditional infrastructure metrics (CPU utilization, HTTP error rates) remained green while model CTR dropped because global metrics masked localized subpopulation degradation.

Systems lesson: Infrastructure telemetry alone cannot detect statistical model failure. High-throughput recommendation systems require cohort-level subpopulation tracking and counterfactual evaluation to detect localized accuracy drift.

14.5.3.1 Data quality monitoring

Model and infrastructure monitoring tracks outputs. By the time output metrics degrade, however, the underlying problem may have existed for days or weeks. Data quality monitoring can expose upstream causes earlier, while output and outcome monitoring reveal whether those changes affect users. The two views are complementary rather than interchangeable. The first input guardrail is executable validation, as listing 14.4 shows: schema expectations reject malformed batches before inference, turning a data-quality assumption into a testable contract.

Listing 14.4: Input Data Validation: Schema validation rules check column existence, data types, null values, and statistical bounds to catch data quality issues before they propagate to model inference.

```
schema.require_column("user_id")
schema.require_type("timestamp", "datetime")
schema.require_non_null("feature_a")

schema.require_range("age", min_value=0, max_value=120)
schema.require_mean_between(
    "purchase_amount", min_value=10, max_value=1000
)
```

26 | **Schema validation:** The rules in listing 14.4 prevent silent data contract violations, such as a feature column changing from an integer to a float. Without this input-level guardrail, downstream model monitoring cannot distinguish a data quality error from a true performance regression, masking the root cause. A schema mismatch in a critical feature can invalidate an otherwise well-formed prediction batch.

Input data validation Schema validation²⁶ catches structural problems before they reach the model. The common rule categories are column existence checks, type enforcement, null detection, and statistical bounds.

Feature distribution monitoring Schema validation catches structural corruption (missing columns, wrong types, null values) but cannot detect the subtler failure mode where data arrives in the correct format but from a shifted distribution. A feature representing user age might pass every schema check while its mean silently migrates from 32 to 45 over three months as a marketing campaign attracts an older demographic. This distributional shift degrades model predictions long before any structural anomaly appears. Statistical distance measures quantify this divergence by comparing current feature distributions against training baselines. Table 14.19 specifies representative alert thresholds for three common metrics, with population stability index (PSI) suited for categorical features, KS statistics for continuous distributions, and Jensen-Shannon divergence for comparing full probability distributions with a symmetric, bounded KL-derived measure.

Table 14.19: Feature Distribution Thresholds: Starting points for drift detection, calibrated in practice to each feature’s sensitivity and business impact. PSI thresholds such as 0.1 and 0.25 are common scorecard-monitoring conventions, while KS and JS thresholds must be calibrated to the feature, sample size, and cost of missed drift. Higher thresholds reduce alert fatigue but risk missing gradual drift.

Metric	Alert Threshold	Use Case
PSI	PSI > 0.25	Categorical and binned features
Kolmogorov-Smirnov statistic	KS > 0.1	Continuous feature distributions
Jensen-Shannon divergence	JS > 0.1	Probability distributions

Understanding these thresholds requires looking at the math. The PSI²⁷ quantifies distributional shift by comparing expected (training) vs. actual (serving) frequencies across bins (section B.2.2 develops the mathematical foundations of KL divergence, PSI, and information theory for systems monitoring). Here n is the number of fixed bins, and expected_i and actual_i are aligned, strictly positive bin proportions that sum to one in each distribution. Equation 14.12 formalizes this:

$$\text{PSI} = \sum_{i=1}^n (\text{actual}_i - \text{expected}_i) \times \ln \left(\frac{\text{actual}_i}{\text{expected}_i} \right) \quad (14.12)$$

Production monitors therefore need a documented zero-bin policy, such as smoothing, and must preserve bin boundaries across comparison periods; PSI is a drift signal, not an automatic retraining trigger.

For discrete or binned distributions, KL divergence provides another comparison. Let p be the monitored serving distribution and q the training reference distribution. Equation 14.13 defines the local KL-specific notation $\mathcal{D}_{\text{KL}}$; elsewhere in this book, $\mathcal{D}(P_t \| P_0)$ denotes a generic statistical divergence in the degradation equation:

$$\mathcal{D}_{\text{KL}}(p \| q) = \sum_x p(x) \log \left(\frac{p(x)}{q(x)} \right) \quad (14.13)$$

Because KL divergence is asymmetric and becomes infinite when $p(x) > 0$ where $q(x) = 0$, operational monitors must preserve direction and apply an explicit support policy, such as smoothing previously unseen bins. To see this in practice, consider a recommendation system monitoring user age: a shift from “younger” to “older” demographics might look subtle on a histogram but generates a clear PSI signal, decomposed bin by bin in table 14.20:

Table 14.20: PSI Worked Example: User age distribution drift from training to serving, decomposed across six age bins. The total PSI is 0.029, well below the 0.1 warning threshold, even though several bins shifted by 3 percentage points. Aggregate PSI summarizes movement across bins; action depends on a calibrated threshold, sample size, and business cost.

Age Bin	Training	Serving	Difference	ln(Serving/Training)	Contribution
18–25	15%	12%	-0.03	-0.223	0.0067
26–35	25%	22%	-0.03	-0.128	0.0038
36–45	20%	18%	-0.02	-0.105	0.0021
46–55	18%	20%	+0.02	+0.105	0.0021

²⁷ | **PSI (population stability index):** PSI is widely used in credit-risk scorecard monitoring to compare expected and observed binned populations; Yurdakul and Naranjo analyze its statistical properties (Yurdakul and Naranjo 2020). The common 0.1 and 0.25 bands are useful operational conventions, not universal statistical laws. ML operations adopted PSI because it works on binned categorical or continuous features and provides an interpretable drift score that non-specialists can review.

Age Bin	Training	Serving	Difference	ln(Serving/Training)	Contribution
56–65	12%	15%	+0.03	+0.223	0.0067
66+	10%	13%	+0.03	+0.262	0.0079

Summing the six bin contributions gives us a total PSI of 0.029 (Stable). The training and serving columns are shown as percentages, while the difference column is expressed in proportion units, so a 3 percentage-point shift appears as ± 0.03 . Even though specific bins shifted by 3 percentage points, the aggregate drift is well below the 0.1 warning threshold. Operational action depends on a calibrated threshold, sample size, and business cost.

Data freshness monitoring Feature stores and data pipelines can become stale without triggering obvious errors. Data freshness monitoring catches that failure mode, and listing 14.5 shows a configuration that monitors feature freshness and triggers fallback behavior when data becomes stale.

Listing 14.5: Data Freshness Alert Configuration: This configuration monitors the `user_purchase_history` feature for staleness, alerting operations teams via PagerDuty and Slack and falling back to default values when the feature exceeds the maximum allowed age.

```
# Example freshness alert configuration
feature: user_purchase_history
max_staleness: 6h
alert_channels: [pagerduty, slack]
on_stale:
  action: fallback_to_default
  default_value: []
```

A freshness policy turns staleness from a silent data defect into an explicit fallback; the same detect-and-respond contract must cover every layer of the monitoring stack.



Checkpoint 14.2: The monitoring stack

ML monitoring is layered, not monolithic. The prompts in this checklist test whether each symptom can be traced to the responsible layer.

- ☐ **Infrastructure utilization:** can you distinguish a 0 percent GPU utilization failure from a 100 percent queuing failure?
- ☐ **Traffic throughput:** can you tell whether a sudden QPS drop points to upstream routing failure or client-side demand change?
- ☐ **Data freshness:** can you detect stale features before plausible-looking predictions reflect yesterday's world?
- ☐ **Data drift:** can you use PSI or KL divergence to decide whether production data still resembles the training distribution?
- ☐ **Model accuracy:** can you account for delayed ground truth labels when interpreting accuracy alerts?
- ☐ **Cohort bias:** can you use cohort analysis to catch subgroup failures that aggregate accuracy hides?

The same stack must watch the data sources that feed the ML system: database replication lag, API endpoint availability, and completion status for extract, transform, load jobs. In one representative incident pattern, a recommendation system detects a material shift in `user_lifetime_value` distribution within two days and traces the issue to a database migration that changed aggregation logic. Without data quality monitoring, this kind of issue can degrade recommendations for weeks before accuracy metrics detect the problem.

Monitoring cost model Observability infrastructure incurs costs that scale with monitoring granularity. Understanding these costs enables rational decisions about monitoring depth vs. budget constraints.

Cost components Monitoring costs break down into four categories, as equation 14.14 decomposes:

$$\text{Monitoring Cost} = C_{\text{ingest}} + C_{\text{storage}} + C_{\text{compute}} + C_{\text{alert}} \quad (14.14)$$

The four C_* terms are cost components over the same accounting window: data ingestion, retained storage, query or dashboard compute, and alert-rule evaluation. Separating them matters because each scales with a different control knob.

Table 14.21 provides representative unit-cost assumptions for each component. Translating these unit costs into a concrete budget estimate clarifies the real expense of monitoring even a single production model:

Table 14.21: Monitoring Cost Components: Illustrative scenario unit costs used for the worked example. Costs scale differently across components: metric ingestion scales with cardinality (number of unique metric series), storage scales with retention, and query costs scale with dashboard usage patterns.

Component	Illustrative Unit Cost	Scaling Factor
Metric Ingestion	\$0.10–0.50 per million data points	Number of metrics $\times$ sample rate
Log Storage	\$0.50–2.00 per GB/month	Log verbosity $\times$ retention period
Query Compute	\$0.01–0.05 per query	Dashboard refresh rate $\times$ users
Alert Evaluation	\$0.001–0.01 per evaluation	Number of alert rules $\times$ check frequency

Napkin Math 14.6: Single-model monitoring budget

Problem: What monthly ingestion, storage, and query cost follows from monitoring a single ML node under these assumptions?

Variables:

- one model with 3 deployment variants (production, canary, staging), each emitting 50 metrics
- Metrics sampled every 15 seconds
- Retention requirement: 30 days
- 2 dashboards (model health, infrastructure), 3 team members, five-minute refresh

Metric ingestion:

- Data points per month: $3 \times 50 \times (4 \text{ samples/min} \times 60 \times 24 \times 30 \text{ days}) = 25.9\text{M}$
- Cost at \$0.30/million: \$7.8/month

Storage:

- At 8 bytes/point compressed: $25.9\text{M} \times 8 \text{ bytes} = 0.2 \text{ GB}$
- Cost at \$1/GB: \$0.21/month

Query compute:

- Queries per month: $2 \text{ dashboards} \times 3 \text{ users} \times (12 \text{ queries/hour} \times 8 \text{ hours/day} \times 22 \text{ days}) = 12,672 \text{ queries/month}$
- Cost at \$0.02/query: \$253.4/month

Total: ~\$261.4/month for a single ML node. Alert evaluation and platform overhead are excluded.

Systems insight: Under these assumptions, cost scales linearly with identical nodes; at platform scale, query-cost optimization becomes increasingly important.

Cost optimization strategies The dominant cost driver in monitoring infrastructure is metric cardinality: high-cardinality labels such as `user_id` or `request_id` create a combinatorial explosion in storage requirements that can dwarf compute costs. Addressing cardinality through sampling or aggregation for high-cardinality dimensions typically yields the largest immediate savings. The second major cost driver is temporal resolution: storing all metrics at 15-second granularity for 30 days is rarely necessary, yet it is the default in many monitoring systems. A tiered retention policy (high-resolution for recent incidents, downsampled data for longer history) preserves debugging fidelity while reducing storage. Dashboard query costs accumulate more subtly: each refresh triggers queries against the metrics backend, and default auto-refresh intervals across dozens of dashboards and users generate continuous query load even when no one is actively watching. Setting slower refresh intervals for noncritical dashboards and auto-pausing inactive tabs can reduce query costs. Finally, alert configuration affects both compute costs and operational effectiveness: consolidating related alerts into multi-condition rules reduces evaluation overhead while also reducing alert fatigue, aligning cost optimization with operational quality.

Cost-benefit framework Justify monitoring investments against incident costs using the monitoring benefit-cost ratio in equation 14.15:

$$\text{Monitoring Benefit/Cost} = \frac{\text{Incidents Prevented} \times \text{Avg Incident Cost}}{\text{Annual Monitoring Cost}} \quad (14.15)$$

If average incident costs \$50,000 (downtime + engineering time + reputation) and monitoring prevents 5 incidents annually at \$50,000 monitoring cost:

$$\text{Benefit/Cost} = \frac{5 \times \$50,000}{\$50,000} = 5 \times$$

This framework helps justify monitoring investments and prioritize which metrics deserve fine-grained observation vs. coarse sampling. The monitoring systems themselves require resilience planning to prevent operational blind spots. When primary monitoring infrastructure fails (Prometheus experiencing downtime or Grafana becoming unavailable), teams risk operating blind during critical periods. Production-grade MLOps implementations therefore maintain redundant monitoring pathways: secondary metric collectors that activate during primary system failures, local logging that persists when centralized systems fail, and heartbeat checks that detect monitoring system outages.

Some organizations implement cross-monitoring where separate infrastructure monitors the monitoring systems themselves, ensuring that observation failures trigger immediate alerts through alternative channels such as PagerDuty or direct notifications. This defense-in-depth approach prevents the catastrophic scenario where both models and their monitoring systems fail simultaneously without detection. A circuit breaker²⁸ adds a further safeguard, automatically routing traffic away from a failing service when its error rate exceeds a threshold. Coordinating these safeguards across many replicated services, with consensus-based alerting and cross-region metric aggregation, is a fleet-scale concern that arises once a model is replicated across regions, beyond the single-node scope here.

14.5.4 Incident response and operational practices

Monitoring and drift detection identify problems; the practices in this section resolve them and sustain operational health over time. Incident response, debugging, and on-call rotations form the human side of the production-monitoring interface, ensuring that statistical signals translate into timely engineering action.

14.5.4.1 Incident response for ML systems

At 2:00 AM, an on-call engineer receives an alert: recommendation click-through rate has dropped 12 percent over the past hour. There is no stack trace, no error log, no crashed process, just a statistical signal that something has changed. The responder must distinguish among four candidate root causes: an upstream data-pipeline failure, model drift, a seasonal traffic pattern, or statistical noise. This ambiguity is common in ML incidents: symptoms may appear as degradation in outcome or

²⁸ | **Circuit breaker pattern:** Automatic failure detection that “opens” when error rates exceed configured thresholds, routing traffic away from failing services. In ML systems, the pattern requires a critical adaptation: prediction accuracy degradation demands different thresholds than service availability failures, because a model returning plausible but incorrect predictions triggers no error signal, leaving the circuit breaker blind to the most dangerous failure mode.

predictive-quality metrics rather than explicit errors, so incident response must account for statistical uncertainty.

Severity classification provides the foundation for prioritizing response in this ambiguous landscape. Table 14.22 defines four priority levels with associated response times, from P0 complete failures requiring 15-minute response to P3 minor anomalies allowing 24-hour investigation.

Once severity is assigned, the incident response process follows a structured checklist whose order narrows the search at each step:

1. Detection determines which monitoring signal triggered the alert.
2. Impact assessment quantifies what percentage of traffic is affected.
3. Responders review recent changes to identify whether any models, features, or data pipelines were deployed.
4. Mitigation options are evaluated, including rollback, fallback enablement, or traffic reduction.
5. Root cause analysis determines whether the issue stems from the model, data, or infrastructure.

Table 14.22: Incident Severity Classification for ML Systems: Response times reflect the urgency and potential business impact of each severity level.

Level	Criteria	Response Time	Example
P0	Complete model failure, serving errors	15 minutes	Model returns null predictions
P1	Significant accuracy degradation (>10%)	1 hour	Recommendation CTR drops 15%
P2	Moderate drift, localized impact	4 hours	One feature shows PSI > 0.3
P3	Minor anomalies, no user impact	24 hours	Training pipeline delay


For P0 and P1 incidents, postmortem documentation is required. These postmortems must include timeline, root cause, user impact, and preventive measures. ML-specific elements include identifying which monitoring gap allowed the issue to reach production and what validation would have caught it earlier.

14.5.4.2 Model debugging: From detection to diagnosis

Incident response triages and mitigates; model debugging identifies root causes. Monitoring detects that something is wrong; debugging determines *why*. ML debugging must account for probabilistic and data-dependent failures in addition to conventional software defects. An incorrect prediction does not throw an exception or generate a stack trace, making systematic debugging essential for resolving ML incidents efficiently.

The debugging decision tree When model performance degrades, work through these diagnostic questions in order. Section A.8 supplies a systematic diagnostic matrix that maps symptoms to D·A·M (Data · Algorithm · Machine) axes.

1. **Is it the data?** Check for upstream data pipeline failures, schema changes, missing values, or distribution shifts. Data is often the first place to look because many production ML failures originate in changing inputs, labels, or feature pipelines.
2. **Is it training-serving skew?** Compare feature values or preprocessing outputs for matched examples across training and serving. KS or PSI can screen for distributional divergence, but cannot establish its cause.
3. **Is it a specific subpopulation?** Slice performance by key dimensions (geography, device type, user segment). Degradation localized to one slice suggests a data coverage or labeling issue.
4. **Is it temporal?** Plot performance over time. Sudden drops prioritize checks for recent deployments and data failures; gradual decline can be consistent with drift, but timing alone does not establish the cause.
5. **Is it the model?** After checking data and pipeline causes, examine model behavior through prediction analysis and feature attribution.



Production debugging starts on the data axis of D·A·M.

Slice analysis This sequence makes slice analysis the first deepening step once a global degradation has been detected, because it tests whether the apparent system-wide problem is actually concentrated in a subpopulation. Performance metrics aggregated across all traffic can mask significant problems in subpopulations: **Slice analysis** exposes that masking, and table 14.23 illustrates how overall accuracy can hide severe degradation in specific segments.

Table 14.23: Slice Analysis Example: Overall accuracy of 91 percent appears acceptable, but tablet users (5 percent of traffic) experience 62 percent accuracy, a severe degradation masked by aggregation. Effective debugging requires systematic slice analysis across key dimensions.

User Segment	Traffic %	Accuracy	Impact
Desktop users	45%	94%	Nominal
Mobile (iOS)	30%	92%	Nominal
Mobile (Android)	20%	88%	Minor degradation
Tablet users	5%	62%	Severe—investigate
Overall	100%	91%	Masks tablet problem

Feature attribution for debugging When slice analysis identifies a problematic segment, feature attribution techniques help identify which features the model relies on within incorrect predictions. Listing 14.6 demonstrates a workflow that uses SHAP values, feature-attribution scores that estimate how much each input feature contributed to an individual prediction, to analyze mispredictions within a specific slice.

Listing 14.6: SHAP-Based Debugging Workflow: This code filters mispredicted tablet examples, computes SHAP values for their model outputs, and plots the selected examples’ feature attributions.

```
# SHAP-based debugging workflow
import shap

# Select mispredicted examples from problematic slice
errors = predictions[
    (predictions.actual != predictions.predicted)
    & (predictions.device_type == "tablet")
]

# Compute SHAP values for error cases
explainer = shap.Explainer(model)
shap_values = explainer(errors[feature_columns])

# Plot feature attributions for the selected errors
shap.summary_plot(shap_values, errors[feature_columns])
```

The attribution plot identifies features the model relied on within the failing slice; it does not establish whether those features caused the errors. Diagnosing stale features, missing coverage, or semantic shift requires checking feature lineage and production distributions.

 Systems Perspective 14.3: Zombie features

Long-lived production models often accumulate features whose original owners, semantics, or intended use have faded. A feature can be deprecated in application code while still flowing through a feature store, training dataset, or serialized model input contract. The model may learn to ignore it, split signal across a duplicate, or depend on a preprocessing artifact that nobody still owns. Removing such a feature is therefore no longer a local cleanup: it becomes a compatibility change that can affect retraining, serving, monitoring, and downstream analysis. Features in ML systems do not disappear just because code owners stop thinking about them. Unused feature payload volume D_{vol} inflates serialization overhead and consumes memory bandwidth BW without improving accuracy. Without explicit deprecation policies and feature-

store governance, models accumulate “dead code” that degrades D_{vol}/BW and complicates debugging (Sculley et al. 2015).

Zombie features show that attribution can reveal what a model should no longer depend on. For individual mispredictions, counterfactual analysis adds the complementary view: the minimal change that would flip a single decision.

Counterfactual analysis For individual mispredictions, counterfactual analysis identifies the minimal change that would flip the prediction: if `session_duration` were 45 seconds instead of 12 seconds, the model would predict “engaged” instead of “churned.” This reveals which feature boundaries drive decisions and whether those boundaries make semantic sense. Counterfactuals that require implausible changes (“user age would need to be -5 years”) often indicate feature engineering problems.

These techniques (decision trees, slice analysis, feature attribution, and counterfactuals) form a debugging toolkit. To apply them consistently, teams codify the process.

Debugging checklist Systematic debugging follows a six-phase checklist: reproduce, isolate, bisect, attribute, validate, and prevent. The ordering is deliberate because each phase narrows the search space for the next. Reproduction comes first because an ML failure that cannot be reproduced on held-out data is often data-dependent, an insight that redirects investigation toward the D-A-M taxonomy’s data layer. Once reproduced, isolation identifies the minimal input set that triggers the failure, transforming a diffuse “the model is wrong” complaint into a specific, testable condition.

Bisection then exploits version history: if the failure correlates with a recent deployment, comparing model versions pinpoints which change introduced the regression. Feature attribution applies these interpretability techniques to identify which input factors drive the erroneous behavior. Validation closes the causal loop by confirming that the hypothesized root cause, when corrected, actually resolves the failure, distinguishing genuine fixes from coincidental improvements.

The final phase, prevention, converts each resolved incident into a monitoring rule or validation check, systematically closing the gap between detection and recurrence. This cumulative hardening can reduce repeated failure modes over time because each incident strengthens the observability infrastructure.

Debugging ML systems requires both systematic methodology and domain expertise. The most effective debugging often comes from engineers who understand both the model architecture and the business context of the predictions.

14.5.4.3 On-call practices for ML systems

The debugging techniques in section 14.5.4.2 work when an engineer is actively investigating an issue during business hours. Production systems, however, fail at 3:00 AM on weekends, and the person responding may not be the one who built the model. Debugging resolves individual incidents; on-call practices sustain operational health over time by ensuring that someone with appropriate expertise is always available and equipped to respond. On-call rotation for ML systems requires specialized practices beyond traditional software operations because ML incidents often manifest as gradual degradation rather than hard failures. A traditional software engineer responding to an alert can typically trace a stack trace to a root cause within minutes. An ML engineer facing a 3 percent accuracy drop must first determine whether the change represents statistical noise, legitimate concept drift, or a critical failure requiring immediate rollback. This distinction demands statistical context rather than simple log analysis.

This ambiguity compounds with delayed impact visibility. Unlike latency spikes that surface immediately in dashboards, ML degradation may take hours or days to manifest in business metrics. A recommendation model that began serving slightly worse suggestions on Monday might not produce measurable revenue impact until Friday, by which time the window for easy diagnosis has closed. Cross-system dependencies further complicate response: ML issues often originate in upstream data systems owned by different teams, requiring coordination across organizational boundaries during incident response. The deepest challenge is that effective response demands understanding model behavior, not infrastructure health alone. A database administrator can restart a crashed service without understanding its business logic, but an ML engineer cannot meaningfully

debug accuracy degradation without understanding the model’s feature dependencies and expected behavior patterns.

These challenges motivate tiered escalation structures that match expertise to incident complexity. Table 14.24 illustrates one possible structure, where primary responders handle routine issues using standardized runbooks while escalation paths connect to specialists capable of deeper investigation. A parallel data on-call role can reduce time to resolution when the root cause lies in an upstream data system.

Table 14.24: ML On-Call Structure: Tiered escalation with parallel data on-call enables efficient incident response. Tier 1 handles routine issues using runbooks; Tier 2 addresses complex debugging; Tier 3 manages critical incidents requiring architectural decisions.

Tier	Responder	Responsibility
Tier 1 (Primary)	ML Engineer	Initial triage, standard runbooks, escalation decisions
Tier 2 (Escalation)	Senior ML Engineer/Data Scientist	Complex debugging, cross-system investigation, model-specific issues
Tier 3 (Critical)	ML Platform Lead	Architecture decisions, major incidents, vendor escalation
Data On-Call (Parallel)	Data Engineer	Data pipeline issues, feature store problems, upstream dependencies

Effective on-call depends heavily on runbook quality. Every production ML model should have documentation covering the model’s purpose, ownership, and business criticality alongside its normal operating parameters: expected latency, throughput, and accuracy ranges that define healthy behavior. Historical incidents and their resolutions provide templates for common failure patterns, while diagnostic commands enable rapid health assessment: how to check recent predictions, feature distributions, and model confidence scores. Critically, runbooks must specify escalation criteria (when to wake up Tier 2 vs. when to rollback without approval) and rollback procedures with step-by-step instructions and expected recovery times. Runbooks written during calm periods save critical minutes during 3:00 AM incidents.

Even well-designed monitoring can generate excessive alerts that erode on-call effectiveness. Alert fatigue, the tendency to ignore or dismiss alerts after experiencing too many false positives, represents a significant operational risk. Teams combat fatigue through consolidation, grouping related alerts so that multiple features drifting simultaneously generate a single notification rather than dozens. Adaptive thresholds that account for weekly and seasonal patterns prevent predictable variations from triggering unnecessary pages. Alerts that are rarely actionable should be retired or recalibrated. When temporary silencing is necessary, accountability mechanisms (requiring a follow-up ticket before snoozing) prevent alerts from being permanently ignored.

Shift handoffs represent another critical practice that distinguishes mature operations. Incoming on-call engineers need context about active incidents and their current status, recent deployments that might cause delayed issues, upcoming scheduled changes such as data migrations or model updates, and any alerts that were suppressed along with the reasoning. Without structured handoffs, context is lost between shifts, and incoming engineers waste time rediscovering information their predecessors already gathered.

Sustainable on-call practices must also address burnout. ML on-call carries particular stress due to incident ambiguity: the uncertainty of not knowing whether an alert represents a real problem demands constant vigilance. Organizations mitigate burnout by limiting consecutive on-call days, providing compensatory time off after high-severity incidents, conducting regular rotation reviews to balance load across team members, and investing in automation that reduces toil. The goal is to make on-call rotations sustainable over years of operation, not to staff them as an afterthought.

Technical monitoring capabilities alone do not ensure operational success. The most sophisticated dashboards fail if no one is responsible for acting on alerts, and the most detailed runbooks languish if team structures do not support their use. Production ML operations require organizational infrastructure paralleling the technical: clear governance, defined roles, and communication patterns that enable cross-functional coordination.

14.5.5 Governance and team coordination

On-call practices address operational emergencies, but production ML also requires proactive governance and cross-functional collaboration. Governance encompasses the policies and practices ensuring that ML models operate transparently, fairly, and in compliance with ethical and regulatory standards. Without it, deployed models may produce biased or opaque decisions, creating legal, reputational, and societal risks. Governance focuses on three core objectives: transparency (interpretable, auditable models), fairness (equitable treatment across user groups), and compliance (alignment with legal and organizational policies). The specific interpretability methods, fairness metrics, and bias detection techniques that operationalize these objectives are examined in chapter 15; MLOps provides the infrastructure to enforce these checks continuously throughout the deployment lifecycle.

Like conventional software compliance, ML governance must span development, deployment, and operation. During development, teams must document model assumptions and training data provenance. At deployment, prerelease audits evaluate fairness and robustness. Postdeployment, the monitoring systems discussed in the previous section must track not only performance degradation but also fairness drift, in which performance or outcome disparities change across user subgroups. Governance policies encoded into automated pipelines ensure that these checks are applied consistently rather than relying on ad hoc human review.

Concretely, a model registry promotion gate might require a signed feature contract, a recorded training-data lineage hash, subgroup metrics above policy thresholds, a canary SLO with rollback criteria, and a named artifact owner before the model can move from staging to production. That gate turns governance from a meeting into a release invariant enforced by the same CI/CD machinery that deploys the model.

Governance establishes policies, but cross-functional collaboration implements them. Machine learning systems are developed and maintained by multidisciplinary teams, and the boundaries between roles create the most failure-prone points in the entire lifecycle. Shared experiment tracking, model registries, and standardized documentation provide the connective tissue that enables reproducibility and eases handoff between specialists. Equally important is shared understanding of data semantics: glossaries, schema references, and lineage documentation ensure that all stakeholders interpret features, labels, and statistics consistently.

Titles and boundaries vary across organizations, but one common decomposition uses the five roles in table 14.25. The table maps each role to a primary responsibility without implying that every organization needs five separate positions:

Table 14.25: ML Team Roles Matrix: Clear role boundaries prevent gaps and overlaps. Data Scientists focus on model quality while ML Engineers handle productionization. Data Engineers own data pipelines while Platform Engineers own MLOps tooling. SREs ensure overall system reliability.

Role	Primary Focus	Key Deliverables	Collaboration Points
Data Scientist	Model development, experimentation, algorithm selection	Trained models, experiment results, performance benchmarks	Hands off to ML Engineer for productionization
ML Engineer	Production ML systems, training pipelines, serving infrastructure	Deployed models, training pipelines, serving systems	Receives from Data Scientist; works with Platform Engineer on infrastructure
Data Engineer	Data pipelines, feature engineering, data quality	Feature pipelines, data quality systems, feature stores	Provides data to Data Scientist; maintains feature store for ML Engineer
Platform Engineer	MLOps infrastructure, tooling, automation	CI/CD pipelines, monitoring systems, compute infrastructure	Enables ML Engineer; maintains shared infrastructure
De-vOps/SRE	Reliability, incident response, system health	SLOs/SLAs, on-call procedures, runbooks	Supports all roles; owns production health

Clear role definitions matter most at handoff points, where work transitions between specialists. The most failure-prone handoff occurs between Data Scientists and ML Engineers: a model that performs well in a Jupyter notebook may fail in production due to undocumented preprocessing

steps, hardcoded file paths, or environment dependencies. Similarly, the handoff from ML Engineers to SREs (Beyer et al. 2016) requires verified monitoring dashboards, configured alerting rules, documented runbooks, and tested rollback procedures. Data Engineers hand off to the broader ML team through feature contracts, formal specifications of schema, freshness SLOs, and quality guarantees that prevent silent pipeline changes from surfacing as mysterious model degradation weeks later. Organizations mitigate these handoff risks through standardized model interfaces, required documentation, and reproducibility requirements that must be verified before each transition.

14.5.5.1 Stakeholder communication

Effective MLOps extends beyond internal team coordination to the broader communication challenges that arise when technical teams interface with business stakeholders. Cross-functional collaboration addresses coordination within technical teams; stakeholder communication bridges technical and business domains. Effective MLOps bridges these domains by translating machine learning realities into terms stakeholders can act on. Machine learning systems add probabilistic performance, data dependencies, and degradation patterns that stakeholders often find counterintuitive.

The most common communication challenge emerges from oversimplified improvement requests. Product managers frequently propose “make the model more accurate” without understanding underlying trade-offs. Effective communication reframes such requests by presenting concrete options: improving accuracy from 85 percent to 87 percent might require substantially more labeled data and a slower model that violates the latency budget. Articulating specific constraints transforms vague requests into informed business decisions.

Translating technical metrics into business impact requires consistent frameworks connecting model performance to operational outcomes. A 5 percent accuracy improvement appears modest in isolation, but contextualizing this as “reducing false fraud alerts from 1,000 to 800 daily customer friction incidents” provides actionable business context.

Figure 14.9 makes the nonlinear error-cost trade-off visible. In this equal-class-weighted example, the optimal operating point is the classification threshold that minimizes the combined cost index for false positives (blocking a legitimate user) and false negatives (missing fraud).




Figure 14.9: The Business Cost Curve: An equal-class-weighted illustrative cost index vs. classification threshold. Technical metrics like ROC curves hide the economic reality: errors have different costs. Here, a false negative (missed fraud) costs \$500, while a false positive (blocked user) costs \$100. The system flags a transaction as fraud when its score exceeds the threshold, so a lower threshold flags more transactions. Because missing fraud is 5× more costly, the optimal threshold shifts left of center, making the system more aggressive at flagging suspicious transactions and accepting more false positives to minimize expensive misses. With equal costs the optimum would sit at threshold = 0.50; the asymmetry pulls it toward threshold ≈ 0.34. A population expected cost would also weight each term by class prevalence. MLOps tunes thresholds as costs and prevalence change.

Incident communication presents another critical challenge. When a model degrades or requires rollback, stakeholders need the current evidence, affected populations, operational impact, mitigation, and remaining uncertainty. A fluctuation should not be dismissed as normal variation before it

is tested, and drift is not merely planned maintenance when it is already harming outcomes. Regular performance reporting establishes the baseline that makes exceptional behavior easier to explain.

Resource justification requires translating technical requirements into business value. Rather than requesting “eight A100 GPUs for model training,” effective communication frames investments as “infrastructure to reduce experiment cycle time from weeks to days, enabling faster feature iteration.” Timeline estimation must account for realistic proportions: data preparation and deployment integration often dominate the schedule, while model development is only one part of the work.

Consider a fraud detection team implementing model improvements. When stakeholders request enhanced fraud capture, the team responds with a structured proposal: increasing the fraction of fraud dollars captured from 92 percent to 94 percent requires integrating external data sources, extending training duration by two weeks, and accepting 30 percent higher infrastructure costs, but would prevent an estimated \$2 million in annual fraud losses while, under a separate 20 percent reduction assumption, reducing false-positive alerts by 50,000 per month.

Through disciplined stakeholder communication, MLOps practitioners maintain organizational support while establishing realistic expectations about system capabilities. This communication competency is as essential as technical expertise for sustaining successful ML operations.

14.5.6 ML test score

A release-readiness assessment needs a shared inventory of the debt patterns that can make a model unsafe to deploy even when its offline metrics look acceptable. Table 14.26 consolidates the patterns discussed throughout this chapter, providing the reference that the assessment rubric in table 14.27 builds on.

Table 14.26: Technical Debt Patterns: Machine learning systems accumulate distinct forms of technical debt from data dependencies, model interactions, and evolving operational contexts. Primary debt patterns, their causes, symptoms, and recommended mitigation strategies guide practitioners in recognizing and addressing these challenges systematically.

Debt Pattern	Primary Cause	Key Symptoms	Mitigation Strategies
Boundary Erosion	Tightly coupled components, unclear interfaces	Changes cascade unpredictably, CACE principle violations	Enforce modular interfaces, design for encapsulation
Correction Cascades	Sequential model dependencies, inherited assumptions	Upstream fixes break downstream systems, escalating revisions	Careful reuse vs. redesign trade-offs, clear versioning
Undeclared Consumers	Informal output sharing, untracked dependencies	Silent breakage from model updates, hidden feedback loops	Strict access controls, formal interface contracts, usage monitoring
Data Dependency Debt	Unstable or underutilized data inputs	Model failures from data changes, brittle feature pipelines	Data versioning, lineage tracking, leave-one-out analysis
Feedback Loops	Model outputs influence future training data	Self-reinforcing behavior, hidden performance degradation	Cohort-based monitoring, canary deployments, architectural isolation
Pipeline Debt	Ad hoc workflows, lack of standard interfaces	Fragile execution, duplication, maintenance burden	Modular design, workflow orchestration tools, shared libraries
Configuration Debt	Fragmented settings, poor versioning	Irreproducible results, silent failures, tuning opacity	Version control, validation, structured formats, automation
Prototype Debt	Rapid prototyping shortcuts, tight code-logic coupling	Inflexibility as systems scale, difficult team collaboration	Flexible foundations, intentional debt tracking, planned refactoring

With those debt patterns in one place, awareness alone is insufficient; teams need a systematic technical debt assessment rubric that transforms subjective “is this system ready?” conversations into quantifiable evaluations. The ML Test Score (Breck et al. 2017) provides a systematic rubric for evaluating production readiness across four categories: data tests, model tests, ML infrastructure tests, and monitoring tests. The paper defines 28 tests in total, seven per section, with partial or full credit for each test. Readiness is tracked by section rather than by a simple grand-total maturity band: a system with strong model tests but weak monitoring still carries production risk. Table 14.27 summarizes representative tests practitioners should implement:

- **Data section:** Validates feature expectations, privacy controls, and whether each feature is beneficial relative to its operational cost.
- **Model section:** Validates reviewed model specifications, hyperparameter discipline, staleness limits, and offline-online metric alignment.
- **Infrastructure section:** Validates reproducible training, rollback, training-serving consistency, and deployment gates.
- **Monitoring section:** Validates alerts for dependency changes, data invariants, skew, and model staleness.

Table 14.27: ML Test Score Checklist: Representative production-readiness tests from the ML Test Score rubric. The original rubric contains 28 tests grouped into four sections of seven tests each; section scores expose whether production risk comes from data validation, model validation, infrastructure, or monitoring rather than hiding the weakness in a single total. Based on Breck et al. (2017).

Category	Test	Implementation Example
Data Tests	Feature expectations are captured in schema	Great Expectations, TFX Data Validation
	All features are beneficial (no unused features)	Feature importance analysis, ablation studies
	No feature's cost exceeds its benefit	Latency/accuracy trade-off analysis
	Data pipeline has appropriate privacy controls	PII detection, access logging
Model Tests	Model spec is reviewed and checked into version control	Git-tracked model configs
	Offline and online metrics are correlated	A/B test validation of offline improvements
	All hyperparameters are tuned	Automated HPO with tracked results
	Model staleness is measured and bounded	Performance decay monitoring
Infrastructure Tests	Training is reproducible	Fixed seeds, versioned data, locked dependencies
	Model can be rolled back to previous version	Model registry with versioning
	Training and serving code paths are tested for consistency	Feature store integration tests
	Model quality is validated before serving	Automated validation gates in CI/CD
Monitoring Tests	Dependency changes result in alerts	Data schema monitoring
	Data invariants hold in training and serving	Distribution comparison tests
	Training and serving features are not skewed	Training-serving skew detection
	Model staleness triggers retraining	Automated retraining pipelines

Quarterly audits against this rubric, prioritizing tests that address the most frequent incident types, reveal where operational investments will yield the highest reliability gains. Checking boxes is necessary but not sufficient. Production readiness requires understanding how practices integrate into a coherent system and how organizations evolve their capabilities over time.

14.6 Design and Maturity Framework

An organization deploying its initial ML model might rely on a hand-run Jupyter notebook, a scheduled cron job, and minimal monitoring. A mature enterprise runs thousands of models through automated pipelines with drift detection, canary deployments, and continuous validation. Both are doing “MLOps,” yet the gap between them spans orders of magnitude in reliability, cost efficiency, and engineering velocity. Deployment case studies show that practical challenges appear across the ML deployment workflow (Paleyes et al. 2022). This chapter uses operational maturity as a systems lens for that progression: organizations evolve from ad hoc experimentation toward fully automated operations, and understanding where a team stands on this continuum is as important as knowing the technical components themselves. Identifying what investments yield the highest returns at each stage guides resource allocation more effectively than adopting tools indiscriminately.

14.6.1 Maturity levels

The ML Test Score assesses individual practices. **Operational maturity** captures something broader: the systemic integration of those practices into a coherent whole. The key distinction is not *which tools* a team has adopted but *how well* infrastructure, automation, monitoring, governance, and collaboration work together across the ML lifecycle. Lifecycle tools such as MLflow address parts of

that workflow (Zaharia et al. 2018), but maturity is the organizational ability to make the pieces work together. Although operational maturity exists on a continuum, distinguishing broad maturity levels helps illustrate how ML systems evolve from research prototypes to production-grade infrastructure.

At the lowest level, ML workflows are ad hoc: experiments run manually, models train on local machines, and deployment involves hand-crafted scripts. As maturity increases, workflows become structured: teams adopt version control, automated training pipelines, and centralized model storage. At the highest levels, systems are fully integrated with infrastructure-as-code, continuous delivery pipelines, and automated monitoring that support large-scale deployment and rapid experimentation.

The distinguishing marker at each stage is not which tools a team adopts but how tightly infrastructure, automation, and monitoring integrate across the lifecycle. Table 14.28 shows that the leap from ad hoc to scalable is primarily an architectural shift from isolated scripts to a cohesive system.

Table 14.28: Maturity Progression: Machine learning operational practices evolve from manual, fragile workflows toward fully integrated, automated systems, impacting reproducibility and scalability. Key characteristics and outcomes at each maturity level emphasize architectural cohesion and lifecycle integration for building maintainable learning systems.

Maturity Level	System Characteristics	Typical Outcomes
Ad Hoc	Manual data processing, local training, no version control, unclear ownership	Fragile workflows, difficult to reproduce or debug
Repeatable	Automated training pipelines, basic CI/CD, centralized model storage, some monitoring	Improved reproducibility, limited scalability
Scalable	Fully automated workflows, integrated observability, infrastructure-as-code, governance	High reliability, rapid iteration, production-grade ML

Consider how a fraud detection system evolves across these maturity levels:

- **Ad hoc:** A data scientist trains a model in a Jupyter notebook, exports it as a pickle file, and hands it to an engineer who deploys it to a single server. When accuracy drops, the data scientist retrains manually by running the notebook again with fresh data. Debugging requires the original data scientist because no one else understands the preprocessing steps.
- **Repeatable:** The training script is version-controlled, with a scheduled Jenkins job that retrains monthly. Features are computed in a SQL script that engineering maintains separately. The model is deployed via container, with basic accuracy monitoring. When the feature SQL changes, the data scientist must manually verify the model still works.
- **Scalable:** Training and serving use the same feature store, reducing skew. A CI/CD pipeline investigates drift when PSI exceeds 0.2, retrains when evidence supports it, validates the new model against the baseline, and deploys via canary release. Monitoring tracks per-merchant accuracy, triggering investigation when specific segments degrade. The entire lineage from raw data to production prediction is auditable.

The investment required to move between levels is substantial and often spans months of engineering effort, but the reduction in incident frequency and debugging time can justify the cost for production-critical systems.

These maturity levels provide a systems lens through which to evaluate ML operations, not in terms of specific tools adopted, but in how reliably and cohesively a system supports the full machine learning lifecycle. Understanding this progression prepares practitioners to identify design bottlenecks and prioritize investments that support long-term system sustainability.

14.6.2 System design implications

Maturity levels describe organizational stages; system design implications describe the architectural consequences. At each level, the system architecture evolves in response to new expectations around modularity, automation, monitoring, and fault tolerance.

In low-maturity environments, ML systems are monolithic: data processing logic embedded in model code, configurations managed informally, and deployments handled through ad hoc scripts. These architectures enable rapid experimentation but lack the separation of concerns needed

for maintainability or safe iteration. As maturity increases, modular abstractions emerge: feature engineering decouples from model logic, pipelines become declarative, and system boundaries are enforced through APIs. At high maturity, ML systems exhibit properties of production-grade software (stateless services, contract-driven interfaces, environment isolation, and observable execution) where data, models, and infrastructure co-evolve through closed feedback loops.

Figure 14.10 captures this architectural reality as an iceberg. What stakeholders see (uptime, the visible tip) represents only a fraction of what must work correctly beneath the surface. The hidden mass below the waterline shows the threats that can sink a system even when it appears healthy: data drift, concept drift, broken pipelines, schema changes, model bias, and underperforming segments. Operational maturity must address all three domains (data health, model health, service health) simultaneously.




Figure 14.10: Uptime Dependency Stack: The waterline separates visible service uptime from data- and model-health failures that can remain hidden while the service stays available; the surrounding labels organize the monitoring surface into data, model, and service health.

The three threat categories in the iceberg map to distinct failure mechanisms. Data health threats (drift, staleness, and schema changes) erode the statistical assumptions a model was trained on, often without any change to the model itself. Model health threats (accuracy degradation, bias amplification, and feedback loops) compound silently because the model continues to produce outputs that appear well-formed even as their quality decays. Service health threats (configuration sprawl, pipeline fragmentation, and stale dependencies) undermine reproducibility and recoverability. Several of these failures can leave the service available, which is why availability monitoring alone cannot cover the full operational state.

14.6.3 Design patterns and anti-patterns

The most sophisticated infrastructure fails without the organizational patterns to operate it effectively. A feature store cannot prevent training-serving skew if no one owns the feature definitions; automated monitoring cannot catch drift if alerts route to the wrong team. As ML systems grow in complexity, organizational patterns must evolve to match.

In mature environments, organizational design emphasizes clear ownership and interface discipline. Platform teams may take responsibility for shared infrastructure and CI/CD pipelines while domain teams focus on model development and business alignment. Interfaces between teams (feature definitions, data schemas, and deployment targets) are well-defined and versioned.

One effective pattern is a centralized MLOps team providing shared services to multiple model development groups. Such structures promote consistency and reduce duplicated effort. Alternatively, some organizations adopt a federated model, embedding MLOps engineers within product teams while maintaining a central architectural function for system-wide integration.

Anti-patterns emerge when responsibilities are fragmented. The tool-first approach (adopting infrastructure tools without first defining processes and roles) results in fragile pipelines and

unclear handoffs. Siloed experimentation, where data scientists operate in isolation from production engineers, leads to models that are difficult to deploy or retrain effectively.

Organizational drift presents another challenge. As teams scale, undocumented workflows become entrenched and coordination costs increase. Organizational maturity must co-evolve with system complexity through communication patterns, role definitions, and accountability structures that reinforce modularity, automation, and observability.

These organizational patterns must be supported by technical architectures handling the reliability challenges of ML systems. MLOps inherits distributed systems challenges but adds complications through learning components requiring adaptations for probabilistic behavior. Conventional services can also return incorrect results without crashing; ML systems add statistical and data-dependent failure modes that availability checks cannot detect.

Circuit breaker patterns must account for model-specific failure modes, where prediction accuracy degradation requires different thresholds than service availability failures. Bulkhead patterns²⁹ become critical when isolating experimental model versions from production traffic. These patterns require resource partitioning strategies that prevent resource exhaustion in one model from affecting others. Ordinary model prediction errors are semantic failures, not Byzantine faults; Byzantine fault tolerance³⁰ provides only a limited analogy for arbitrary component failures.

Consensus algorithms establish agreement among nodes; they cannot determine whether a model prediction is correct when ground truth is delayed or unavailable. These reliability patterns nevertheless help distinguish robust MLOps implementations from fragile ones when their guarantees are applied within scope.

14.6.4 Contextualizing MLOps

Best practices are rarely deployed in pristine environments. Every ML system operates within a specific context that shapes how practices are implemented: physical constraints (edge compute, power budgets), regulatory requirements (healthcare, finance), or organizational realities (team size, skill distribution). A standard CI/CD pipeline may be infeasible without direct host access; monitoring may require indirect signals or on-device anomaly detection; data collection may be limited by privacy regulations. These adaptations are expressions of maturity under constraint, not departures from the principles.

At the highest levels of operational maturity, the single-model practices established here become building blocks for larger organizational capabilities. Organizations operating many ML nodes simultaneously often consolidate into platform architectures that provide shared infrastructure, centralized governance, and economies of scale. The transition from individual ML nodes to platform-scale infrastructure introduces qualitatively different challenges (cross-model resource allocation, system-level observability, fault tolerance for interdependent AI systems) that extend beyond this chapter's single-model scope. Sound ML node practices are prerequisites for platform success because gaps in single-model monitoring, testing, or deployment multiply across the model portfolio.

14.6.5 MLOps investment economics

The operational benefits of MLOps become persuasive only when the investment matches the model's production value. For a single ML node, the decision is whether deployment speed, incident reduction, and monitoring coverage justify the operational spend; for a portfolio, the same economics compound into platform investment.

14.6.5.1 Single-model MLOps investment

For a single production ML system, the first threshold is the annual cost of making the node observable, deployable, and recoverable. Table 14.29 summarizes the main cost categories:

²⁹ | **Bulkhead pattern:** This pattern partitions system resources to contain failures within isolated zones. For isolating experimental models, a bulkhead dedicates a fixed compute and memory budget to the new version. This resource partition ensures that a catastrophic failure in the experiment, such as a memory leak, cannot exhaust all available resources and cause a system-wide production outage.

³⁰ | **Byzantine fault tolerance:** The classic Byzantine model concerns nodes that may behave arbitrarily or send conflicting messages. In the synchronous unauthenticated, or oral-messages, setting, tolerating f Byzantine faults requires at least $3f + 1$ participants (Lamport et al. 1982). ML ensembles are only an analogy because prediction errors can be correlated, and agreement does not establish semantic correctness.

Table 14.29: Single-Model MLOps Investment: Illustrative planning inputs for operationalizing one production ML system. Open-source tooling (MLflow, Feast) can reduce software costs; cloud-managed services trade higher unit costs for reduced engineering overhead.

Component	Illustrative Assumption	Justification
CI/CD pipeline setup	\$10–30K one-time	Reduces deployment time from days to hours
Monitoring and alerting	\$2–10K/year	Catches degradation before user impact
Feature store (basic)	\$5–20K/year	Reduces one source of training-serving skew
Model registry	\$0–5K/year	Enables rollback, audit trails
Engineering time	1–2 FTE-months setup	Initial automation and integration

14.6.5.2 Single-model ROI calculation

The return threshold then depends on model criticality: a revenue-facing model can justify more operational spend because avoided incidents and deployment-time savings have measurable value. Equation 14.16 formalizes that single-node calculation:

$$\text{Annual Benefit/Cost} = \frac{\text{Incidents Avoided} \times \text{Avg Incident Cost} + \text{Time Savings} \times \text{Hourly Cost}}{\text{Annual MLOps Investment}} \quad (14.16)$$

where Incidents Avoided is the count of production failures the tooling prevents per year and Avg Incident Cost is the loss per failure, so their product is the value of avoided incidents; Time Savings is the engineer-hours that automation reclaims per year and Hourly Cost is the loaded labor rate, so their product is the value of recovered labor; the denominator is the annual cost of the tooling itself. The ratio expresses every dollar of investment in dollars returned.

For a model generating \$1M annual revenue with:

- 4 incidents/year avoided (at \$25K each) = \$100K saved
- 20 hours/month deployment time saved (at \$150/hr) = \$36K saved
- MLOps investment of \$30K/year

$$\text{Benefit/Cost} = \frac{\$100K + \$36K}{\$30K} = 4.53 \times$$

14.6.5.3 When to invest more

The 4.53× benefit-cost ratio means the investment is not justified by tooling elegance; it is justified because the model is expensive enough that preventing incidents and shortening deployments outweigh the annual platform spend. The returns from single-model MLOps practices compound when teams add additional models. The transition from operating several independent ML nodes to building a centralized platform involves different economics entirely, including shared infrastructure amortization, platform team overhead, and cross-model coordination costs.

For single-model operations, invest in MLOps in proportion to model criticality. A model driving \$10M in annual revenue justifies more operational rigor than an internal analytics model. Monitoring, lineage, and repeatable deployment are often the first investments. Add a feature store when shared feature definitions and online-offline parity become material problems, and automate retraining only when evidence supports the trigger and the validation loop can contain a bad update.

The preceding technical infrastructure and economic framework provide the foundation; the case studies in section 14.7 demonstrate how these elements combine in production systems. Each case demonstrates specific implementations of the five foundational principles, identifying where reproducibility appears, how observable degradation is achieved, and what triggers automation.

14.7 Case Studies

A battery-powered sleep-tracking ring and AI/ML-based medical software governed by FDA lifecycle expectations (U.S. Food and Drug Administration 2025) face different operational constraints. The principles, patterns, and infrastructure examined throughout this chapter converge differently depending on the deployment context. An Oura-inspired edge design shows how pipeline debt and

configuration management challenge resource-constrained environments; ClinAIOps shows how feedback loops and governance requirements reshape healthcare operations. The Oura lifecycle in section 14.7.1 is hypothetical; the cited study supports offline data and model evaluation, not claims about Oura’s production OTA or MLOps implementation. The comparison starts with the shared principles, because the domains differ most in how those principles are implemented.

Table 14.30 lays out how the two environments could implement the five foundational MLOps principles. Domain constraints (edge hardware, clinical regulation) reshape *how* each principle is realized without changing *which* principles matter. In the Oura-inspired design, polysomnography (PSG) refers to the clinical sleep-study measurements used as reference labels.

The principles stay stable, but their implementation changes with the deployment regime. Edge systems spend the automation budget on battery, telemetry, and constrained updates; clinical systems spend it on auditability, validation gates, and accountable human control. The two case studies that follow trace how each environment earns those entries.

Table 14.30: MLOps Principles by Case Study: Side-by-side mapping of the five foundational MLOps principles to a hypothetical Oura-inspired edge design and the ClinAIOps framework, showing how domain constraints reshape implementation without changing which principles apply.

Principle	Oura-inspired edge design	ClinAIOps
Reproducibility	Versioned synchronized wearable and PSG datasets	Audit trails, decision provenance
Separation of concerns	Independent data, training, and serving layers with edge-specific deployment pipeline	Distinct clinical validation and deployment stages with regulatory compliance isolation
Consistency	PSG-aligned preprocessing across training and on-device inference	Standardized clinical data pipelines ensuring training-serving parity
Observable degradation	On-device anomaly detection, limited telemetry	Cohort-specific monitoring, outcome tracking
Cost-aware automation	Battery-aware retraining triggers, CI/CD for edge balancing accuracy and resource cost	Automated model updates with human-in-the-loop gates balancing update cost and patient risk

14.7.1 Oura-inspired edge design

The Oura Ring provides the offline evidence behind this Oura-inspired MLOps design. The published work supports the clinical data collection and sleep-stage evaluation described later (Altini and Kinnunen 2021), while the versioning, telemetry, over-the-air deployment, and iterative refinement cycle are a hypothetical architecture rather than a reported account of Oura’s production lifecycle. The constraints imposed by a battery-powered ring with limited compute make every MLOps decision visible in a way that cloud-scale systems can obscure.

14.7.1.1 Context and motivation

The Oura Ring is a consumer-grade wearable monitoring sleep, activity, and physiological recovery through embedded sensing and computation. By measuring motion, heart rate, and body temperature, the device estimates sleep stages and delivers personalized feedback. An Oura-inspired design may place selected preprocessing and inference on the device while other processing occurs on a phone or in the cloud.

The central objective was improving sleep stage classification accuracy to align more closely with polysomnography (PSG),³¹ the clinical gold standard. Initial evaluations showed 57 percent four-stage sleep classification accuracy for an accelerometer-only model, compared with 79 percent for models that included autonomic nervous system and circadian features. Published human PSG inter-scorer reliability is about 82 percent to 83 percent, framing the remaining gap between wearable inference and expert clinical scoring. The 22 percentage-point gain closes roughly 84.6–88 percent of the baseline-to-human-agreement gap, although inter-scorer agreement is not a ceiling on accuracy against a fixed or adjudicated reference. This discrepancy prompted an effort to re-evaluate data collection, preprocessing, and model development workflows.

To overcome performance limitations, the Oura team constructed a diverse dataset grounded in clinical standards through a study involving 106 participants from three continents (Altini and Kinnunen 2021). Each participant wore the Oura Ring while simultaneously undergoing PSG,

³¹ | **Polysomnography (PSG):** A multi-parameter sleep study that provides the clinical ground truth data for this classification task. This ‘truth’ is inherently noisy; expert human scorers interpreting the same PSG recordings agree with each other at about 82 percent–83 percent reliability (Altini and Kinnunen 2021). Inter-scorer agreement indicates label uncertainty but does not impose an accuracy ceiling against a fixed or adjudicated reference.

yielding 440 nights of data and 3,444 hours of time-synchronized recordings that aligned wearable sensor data with validated sleep annotations. The scale and diversity of the collection captured physiological variation as well as environmental and behavioral factors critical for generalizing across a real-world user base.

The study consolidated synchronized accelerometer, temperature, heart-rate, heart-rate-variability, and PSG data from research Oura rings, then resolved temporal alignment and preprocessing requirements for downstream model development. A hypothetical production workflow would add robust versioning and lineage tracking to avoid unstable dependencies that commonly plague embedded ML systems.

With high-quality data in place, the study's next question was whether extra sensing improved sleep-stage classification. The researchers evaluated four offline configurations that incrementally added temperature, heart-rate variability, and circadian features to an accelerometer-only baseline. Through 5-fold cross-validation against PSG annotations, the enhanced models achieved 79 percent four-stage classification accuracy, an improvement from the 57 percent accelerometer-only baseline (Altini and Kinnunen 2021). These offline gains motivate the hypothetical production design, but its reproducible training, versioning, conversion validation, and energy controls are inferred requirements rather than results reported by the study.

Following validation, deployment shifted the problem from model quality to update safety. An Oura-like edge deployment must decide which parts of the model run continuously on-device, which richer signals can be used under looser memory and battery budgets, and how model updates reach devices already in the field. To keep that split maintainable, the operational toolchain needs reproducible model conversion, versioned artifacts, and authenticated OTA³² update procedures that preserve consistency across devices in the field.

The operational lesson is that edge MLOps is not governed by accuracy alone; it is governed by accuracy under battery, privacy, telemetry, and weak ground truth constraints. Consider the DS-CNN (Tiny Constraint) archetype from table 14.4, where monitoring relies on operational metrics such as duty cycle and false positive rate rather than continuous ground-truth labels, and an illustrative quarterly cadence governs OTA updates. Operationalizing the transition from 57 percent accelerometer-only accuracy to 79 percent multi-sensor accuracy would require systematic configuration management across data collection, feature sets, model architectures, and deployment targets.

Those constraints explain how the foundational principles appear without repeating them as a checklist. Versioned wearable and PSG datasets make each model traceable to the evidence used to train it. Modular tiered architectures keep data collection, model training, and on-device serving separate enough that quantization, pruning, and fallback policies can change without destabilizing the whole pipeline. PSG-aligned preprocessing preserves consistency between training and on-device inference, while privacy-preserving telemetry makes degradation observable through duty cycle, battery impact, inference failures, confidence, anomaly rates, and periodic labeled studies. OTA deployment then becomes the cost-aware automation boundary: updates must justify their accuracy gain against battery impact, validation burden, and the risk of changing software on a device worn continuously by users.

This case exemplifies how MLOps principles adapt to domain-specific constraints. When machine learning moves into clinical applications, additional complexity emerges, requiring frameworks that address regulatory compliance, patient safety, and clinical decision-making.

14.7.2 ClinAIOPs case study

Healthcare ML deployment presents challenges extending beyond resource constraints. General-purpose MLOps frameworks require additional controls in domains with extensive human oversight, domain-specific evaluation, and ethical governance. Continuous therapeutic monitoring (CTM)³³ exemplifies a domain where MLOps must evolve to meet clinical integration demands.

CTM uses wearable sensors to collect real-time physiological and behavioral data from patients. AI systems must be integrated into clinical workflows, aligned with regulatory requirements, and designed to augment rather than replace human decision-making. General-purpose MLOps practices provide technical lifecycle controls, but they do not by themselves prescribe clinical responsibility, evidence standards, or decision authority.

³² | **Over-the-air (OTA) updates:** The mechanism used to deploy optimized models to devices already in the field, bypassing the need for physical access. The small footprint from quantization and pruning matters because constrained edge networks may need to transmit only changed model artifacts rather than complete application bundles. This process makes consistency a critical concern; a failed update can corrupt the on-device model, breaking the ML pipeline until a future connectivity window allows for a fix.

³³ | **Continuous therapeutic monitoring (CTM):** Healthcare approach using wearable sensors for real-time physiological data collection and personalized treatment adjustments. CTM forces MLOps to confront constraints absent in typical deployments: feedback loops must include human-in-the-loop approval for safety-critical decisions, retraining requires clinician-validated labels rather than implicit signals, and model updates must satisfy regulatory compliance before deployment. These constraints reshape every MLOps principle, making CTM a stress test for operational maturity.

ClinAIOps (Chen et al. 2023), a framework for operationalizing AI in clinical environments, shows how MLOps principles must evolve for regulatory and human-centered requirements. Unlike conventional MLOps, ClinAIOps directly addresses feedback loop challenges by designing them into the system architecture. The framework's structured coordination between patients, clinicians, and AI developers represents practical implementation of governance and collaboration principles.

Clinical environments therefore extend the ordinary operational problem. Healthcare requires coordination among diverse human actors, clinical decisions depend on personalized care and shared accountability, and health data is subject to strict privacy and governance requirements. ClinAIOps presents a framework that balances technical rigor with clinical utility and operational reliability with ethical responsibility.

14.7.2.1 Feedback loops

Three interlocking feedback loops enable safe, adaptive integration of machine learning into clinical practice. Figure 14.11 maps these loops as a circular flow among three stakeholders. Patients contribute continuous monitoring data from wearable sensors and receive bounded AI-assisted guidance. Clinicians receive AI-generated summaries, alerts, and recommendations, then apply clinical judgment by setting therapy regimens and approval limits. AI developers receive continuous feedback from patients and clinicians, using real-world performance and workflow signals to improve models and deployment processes. The outer loop connecting all three stakeholders represents the full governance cycle.




Figure 14.11: ClinAIOps Feedback Loops: The cyclical framework coordinates patients, clinicians, and AI developers to support continuous model improvement and safe clinical integration. Patients and clinicians use AI outputs in care workflows, while AI developers receive feedback from both groups to refine models and operations. Source: (Chen et al. 2023).

Each feedback loop plays a distinct yet interconnected role; together, these loops enable adaptive personalization, maintain clinician control, and promote continuous model improvement based on real-world feedback:

- The **patient treatment loop** captures real-time physiological data and uses bounded AI outputs to support patient self-management.
- The **clinician oversight loop** ensures AI-assisted recommendations are reviewed, limited, and refined under professional supervision.
- The **developer feedback loop** gives AI developers continuous feedback from patients and clinicians so models, interfaces, and monitoring workflows can improve.

Patient treatment loop The patient treatment loop enables personalized therapy optimization through continuous physiological data from wearable devices. Patients wear sensors such as continuous glucose monitors or ECG-enabled wearables that passively capture health signals.

The AI system analyzes these data streams alongside clinical context from electronic medical records, generating individualized recommendations for treatment adjustments. Treatment suggestions are tiered: minor adjustments within clinician-defined safety thresholds may be acted upon directly by the patient, while significant changes require clinician approval. This structure maintains human oversight while enabling high-frequency, data-driven adaptation.

Clinician oversight loop The clinician oversight loop introduces human oversight into AI-assisted decision-making. The AI generates treatment recommendations with interpretable summaries of patient data including longitudinal trends and sensor-derived metrics.

For example, an AI model might recommend reducing antihypertensive medication for a patient with consistently below-target blood pressure. The clinician reviews the recommendation in context and may accept, reject, or modify it, and this feedback refines model alignment with clinical practice. Clinicians also define operational boundaries that ensure only low-risk adjustments are automated, preserving clinical accountability while integrating machine intelligence.

Developer feedback and patient-clinician coordination Developer feedback and patient-clinician coordination shift clinical interactions from routine data collection to higher-level interpretation, shared decision-making, and model improvement. With AI handling data aggregation and trend analysis, clinicians engage more meaningfully: reviewing patterns, contextualizing insights, and setting personalized health goals.

For example, in diabetes management, a clinician may use AI-summarized data to guide discussions on dietary habits and physical activity. Visit frequency adjusts dynamically based on patient progress rather than fixed intervals. This positions the clinician as coach and advisor, interpreting data through the lens of patient preferences and clinical judgment. Feedback from these interactions gives AI developers evidence about model behavior, interface usability, and workflow fit.

14.7.2.2 Hypertension case example

Hypertension management illustrates how the three ClinAIOps loops work in practice. Because it affects a large share of adults and requires individualized, ongoing therapy adjustments, it is an ideal candidate for continuous therapeutic monitoring.

Data infrastructure Research systems estimate systolic blood pressure indirectly from ECG, photoplethysmography (PPG),³⁴ pulse-transit-time, and heart-rate features (Q. Zhang et al. 2017). In a deployed hypertension workflow, those signals may be augmented by accelerometer data for activity context and self-reported medication adherence logs. Accuracy depends on validation, calibration, and regulatory authorization; consumer wrist or ring claims should not be treated as clinically reliable without such evidence. When validated for the intended population and setting, this multimodal data stream, integrated with electronic health records, can form the foundation for personalized AI recommendations.

Loop implementation Figure 14.12 shows two of the three feedback loops across the top and patient-clinician coordination below. The upper-left panel illustrates the patient treatment loop, where the patient monitors blood pressure and receives bounded titration recommendations that the AI system can issue within clinician-defined safety thresholds; significant changes require explicit approval. The upper-right panel depicts the clinician oversight loop, where longitudinal trend summaries flow from the AI system to the clinician, and the clinician sets approval limits and receives alerts for clinical risk events such as persistent hypotension or hypertensive crisis. The lower panel captures the patient-clinician coordination that emerges once routine data collection moves to the AI loop: appointments shift to higher-level discussions of lifestyle factors and shared decision-making. The third loop (developer feedback) is not depicted in the figure; it is described in section 14.7.2.1 as the channel by which real-world workflow signals from both patients and clinicians inform model and interface improvements.

The three panels make the accountability boundary explicit: routine monitoring can be automated only inside clinician-defined limits, while escalation, adverse-event review, and treatment trade-offs remain human responsibilities. That boundary is the point where ordinary MLOps practices need the additional clinical coordination summarized next.

³⁴ | **Photoplethysmography (PPG)**: Optical technique that infers blood-volume changes from variations in detected light after illuminating tissue. For ML operations, PPG introduces a data quality challenge absent in controlled environments: motion artifacts from wrist movement corrupt the signal, creating a data drift pattern where the same physiological state produces different input distributions depending on user activity. Models must either filter corrupted windows before inference or learn to be robust to motion noise, and monitoring must distinguish genuine physiological changes from artifact-induced distribution shift.




Figure 14.12: Hypertension Management Loops: The three panels, labeled patient-AI, clinical-AI, and patient-clinical, each carry a two-way exchange. The AI issues bounded dose titrations within limits the clinician sets and escalates severe readings, which leaves the patient-clinical panel for trend review, adverse-event checks, and treatment decisions. The developer feedback loop of ClinAIops is not among them and is discussed in the prose. Source: (Chen et al. 2023).

14.7.2.3 MLOps vs. ClinAIops comparison

The hypertension case illustrates the ClinAIops-MLOps comparison. General-purpose MLOps provides the technical lifecycle foundation; high-stakes clinical deployment adds explicit constructs for human decision-making, evidence, and accountability.

ClinAIops extends beyond technical infrastructure to support complex sociotechnical systems, embedding machine learning into contexts where clinicians, patients, and stakeholders collaboratively shape treatment decisions. Table 14.31 contrasts these approaches across eight dimensions.

Table 14.31: Clinical AI Operations: General-purpose MLOps supplies technical lifecycle controls; ClinAIops extends them with clinical workflows, evidence, human oversight, and accountability.

	General-purpose MLOps emphasis	ClinAIops extension
Focus	Technical model lifecycle	Human and AI decision-making
Stakeholders	ML, data, platform, and operations teams	Adds patients, clinicians, and clinical governance
Feedback loops	Monitoring, retraining, and release	Adds treatment, clinician, and developer feedback
Objective	Reliable, governed ML operations	Safe, effective, and accountable care support
Processes	Automated pipelines and operational controls	Integrates clinical workflows and human gates
Data considerations	Lineage, quality, freshness, and access control	Adds consent, clinical provenance, and protected data
Model validation	Predictive, operational, and slice-level metrics	Adds clinical utility, safety, and cohort outcomes
Implementation	Technical integration and operational ownership	Adds clinical accountability and stakeholder incentives

The table’s central distinction is that clinical deployment changes who owns the risk. Technical performance remains necessary, but it is not sufficient when a recommendation affects care decisions. The ClinAIops framework therefore changes the governing constraint from device efficiency to clinical accountability. The model participates in care, but it cannot own the clinical decision. Each recommendation and subsequent clinician action must be traceable to the relevant inputs, model and software versions, and output metadata; otherwise the system cannot support audit, review, or outcome analysis. Separation of concerns becomes a safety mechanism rather than only a software design preference: automated data collection from wearables, AI recommendations, clinician diagnosis, treatment decisions, and developer workflow improvement each need explicit boundaries and human gates at critical decision points.

The same accountability requirement changes monitoring. Standardized clinical data pipelines support training-serving parity, but clinical validation also has to compare recommendations against standard-of-care outcomes, prospective evidence, and cohort-specific effects. Observable degradation is measured through blood pressure control, adverse events, clinician overrides, and subgroup outcomes, not just model metrics. Feedback loops are therefore not technical debt in this setting; patient treatment, clinician oversight, and developer feedback loops are intentional mechanisms that improve care while keeping authority with humans. Cost-aware automation operates inside

those gates: updates can be automated only when their expected benefit justifies validation cost and patient risk, and conservative recommendations or uncertainty flags must route low-confidence cases back to clinical review.

14.7.3 Case study synthesis

The Oura-inspired design and ClinAIOPS cases separate stable MLOps principles from the deployment constraints that reshape their implementation. The hypothetical ring lifecycle is a resource-envelope case: the operational system must preserve reproducibility, consistency, and observable degradation while battery, telemetry, and weak ground truth limit what can be measured and updated on the device. ClinAIOPS is an accountability-envelope case: the same principles apply, but validation, audit trails, and human gates dominate because the model influences clinical action.

The shared engineering lesson is that MLOps maturity is not tool accumulation. It is the ability to identify the governing constraint, choose the operational controls that match it, and preserve evidence when the model changes. Production ML systems can fail when teams apply code-focused operational intuitions without accounting for statistical behavior and changing data, which is why the chapter closes by naming the fallacies and pitfalls that these two case studies help expose.

14.8 Fallacies and Pitfalls

These fallacies and pitfalls capture common errors that waste engineering resources, trigger production incidents, and cause silent accuracy degradation. Each connects to specific sections detailing the underlying mechanisms and solutions.

Fallacy: *MLOps is just applying traditional DevOps practices to machine learning models.*

Engineers assume standard CI/CD pipelines transfer directly to ML, but production ML requires specialized infrastructure. As section 14.4.2.1 showed, ML pipelines add data validation, model training, performance evaluation, artifact registration, and deployment gates that make them slower and more stateful than conventional software pipelines. Traditional DevOps can release deterministic services frequently; ML systems without specialized tooling often slow down because retraining and validation are stateful. Standard CI/CD tools do not by themselves handle feature stores, model registries, or drift detection. A recommendation system deployed using conventional DevOps can lose accuracy because the pipeline lacks training-serving consistency checks. Organizations that adopt DevOps without ML adaptations optimize the computational reliability of their infrastructure while neglecting the statistical behavior of their models, encountering silent model degradation, training-serving skew, and data quality failures that evade conventional testing.

Pitfall: *Treating model deployment as a one-time event rather than an ongoing process.*

Teams view deployment as a terminal milestone analogous to shipping software releases, but model quality can change with data drift and distribution shift. Section 14.5.3.1 establishes PSI as one useful distribution-shift signal whose thresholds must be calibrated to the feature and business risk; crossing one should trigger investigation, not automatic retraining. A fraud detection model can move from below the warning threshold to above the review threshold within months, while outcome metrics determine whether initially acceptable accuracy has degraded. The locally fitted cost model in section 14.4.2.2 gives $T^* \approx \sqrt{\frac{2C}{Q \cdot V \cdot \text{Accuracy}_0 \cdot \gamma}}$ under its assumptions. Production ML requires continuous monitoring of feature distributions and performance metrics, with retraining after diagnosis and validation.

Fallacy: *Automated retraining ensures optimal model performance without human oversight.*

Engineers assume automated pipelines handle all maintenance scenarios, yet automation cannot detect all failure modes. Automated retraining can perpetuate biases in corrupted training data, trigger updates during peak traffic, or deploy models that pass aggregate validation but degrade edge cases. A news recommendation system retrained on weekend data might exhibit lower weekday engagement because user behavior differs sharply across weekday vs. weekend contexts. Effective MLOps requires escalation protocols for anomalous validation results, manual approval for unusual metric patterns, and override capabilities when automation produces questionable outcomes.

Pitfall: *Focusing on technical infrastructure while neglecting organizational and process alignment.*

Organizations invest in MLOps platforms expecting tooling to solve deployment problems, but sophisticated infrastructure fails without cultural transformation. MLOps demands coordination between data scientists optimizing for accuracy, engineers prioritizing latency, and business stakeholders focused on impact. A retail company may deploy feature stores and model registries yet maintain a slow deployment cadence because data scientists and engineers operate in isolation. Successful MLOps requires cross-functional teams with unified objectives, shared on-call rotations building empathy across roles, and incentive structures rewarding production reliability alongside model performance.

Fallacy: *Training and serving environments automatically remain consistent once pipelines are established.*

Teams assume that feature computation produces identical values across training and serving after initial pipeline setup, but training-serving skew emerges from subtle inconsistencies in preprocessing logic, timezone handling, or dependency versions. Section 14.4.1.2 demonstrates how a feature store reduces this risk by centralizing feature definitions and comparing feature distributions across environments. An e-commerce ranking model that computes `session_length` using wall-clock time in training but processing time in serving can suffer material accuracy loss that persists until someone compares feature distributions directly. Without centralized feature stores and automated consistency validation, skew detection can take weeks as degradation gradually becomes visible in aggregate metrics.

Pitfall: *Assuming comprehensive monitoring prevents all production incidents.*

Engineers believe sufficient metrics and dashboards eliminate surprise failures, but monitoring creates blind spots when teams track outputs without validating inputs. Section 14.5.3.1 establishes that input validation can detect issues before they degrade predictions, yet many ML systems in practice monitor only accuracy and latency. A recommendation system can track click-through rate while ignoring feature staleness, missing embeddings that are hours out of date due to database replication lag. This can create engagement degradation before outcome monitoring triggers alerts. Systems monitoring only outputs can detect failures late; adding data-quality monitoring can reduce time to detection. Production ML requires layered monitoring with explicit SLOs for data freshness, schema validation, feature distributions, model outputs, and business metrics. Monitoring infrastructure itself needs redundancy to prevent blind operation during platform failures.

Fallacy: *Accuracy is the first production signal to monitor.*

Teams instrument production with accuracy dashboards and assume degradation will appear there first. Accuracy is often a lagging indicator. A model's aggregate accuracy can remain stable even as the input distribution drifts or subgroup behavior shifts. By the time accuracy visibly degrades, the drift may have been accumulating for weeks. Monitoring input distributions with PSI or KL divergence (section 14.5.3.1) can flag change earlier and allow investigation, but labeled outcomes and diagnosis determine whether retraining is appropriate.

Pitfall: *Routing leading-indicator alerts to a different channel than accuracy alerts.*

Teams that do instrument drift and freshness signals often wire them to a dashboard or a low-priority queue separate from the on-call path that handles accuracy regressions, so the early warning fires but no one is paged. A leading indicator only buys time if it reaches an owned response path with severity and runbooks calibrated to its operational impact; it need not page with the same urgency as confirmed accuracy loss. The operational goal is to make accuracy the confirmation signal rather than the first sign of trouble, which holds only when the earlier signals are acted on appropriately.

Together, these failures show that observability is not merely a dashboard problem. Signals must cover data, models, and services, then reach people who own the response; production discipline joins the technical pipeline to the human response path.

14.9 Summary

MLOps exists because machine learning systems add failure modes beyond those of conventional software. A crashed server turns availability dashboards red, but an available service can also return incorrect results. A degrading model adds a statistical failure mode in which predictions continue

while accuracy erodes without an availability signal. This difference explains why conventional operational practices require additional controls for ML and why machine learning operations emerged to close that observability gap.

The five foundational principles introduced at the chapter's opening (section 14.2.1) provide an evaluation framework that applies regardless of scale or domain. Reproducibility through versioning addresses the root cause of many production incidents: untracked artifacts including data versions, configuration changes, and environment drift that make debugging impossible and rollbacks unreliable. Separation of concerns contains the blast radius when changes are required, preventing the boundary erosion and correction cascades that transform local fixes into system-wide regressions. The consistency imperative targets training-serving skew, the silent accuracy killer that appears when feature computation diverges between pipelines; feature stores centralize feature definitions and serving paths, though parity still requires validation. Observable degradation transforms the abstract "silent failure" problem into actionable alerts through layered monitoring that tracks data freshness, feature distributions, model outputs, and business metrics. Cost-aware automation replaces arbitrary retraining schedules with a fitted staleness cost model ($T^* \approx \sqrt{\frac{2C}{Q \cdot V \cdot \text{Accuracy}_0 \cdot \gamma}}$) whose use depends on its assumptions.

The infrastructure components examined throughout the chapter directly implement these principles across the three critical interfaces introduced at the chapter's opening. Feature stores and data versioning address the Data-Model Interface by reducing training-serving inconsistency. CI/CD pipelines and model registries address the Model-Infrastructure Interface by enforcing reproducibility and enabling rollback. Monitoring systems, incident response frameworks, and on-call practices address the Production-Monitoring Interface by making degradation observable and actionable. The retraining decision framework enables cost-aware automation by connecting measured degradation to economic thresholds. The case studies demonstrated that domain constraints reshape *how* principles are implemented without changing *which* principles matter: the Oura-inspired design uses the study's 57 percent to 79 percent offline accuracy improvement to motivate resource-aware validation and update controls rather than claiming a reported production lifecycle. ClinAI Ops showed why high-stakes clinical use makes graceful-degradation and human-oversight controls important, with three feedback loops functioning as architectural patterns rather than operational overhead.



Key Takeaways: Perfectly available, perfectly wrong

- **ML systems can fail silently, and drift metrics signal change:** ML quality can degrade as distributional divergence $\mathcal{D}(P_t \| P_0)$ grows, although divergence alone does not determine accuracy. A model can maintain perfect uptime while accuracy falls. Outcome monitoring is essential, not uptime tracking alone.
- **Training-serving skew can silently change accuracy:** Feature stores reduce skew by centralizing feature definitions and serving paths, though parity still requires validation.
- **Retraining is an engineering optimization, not a guess:** The fitted staleness cost function ($T^* \approx \sqrt{2C/(Q \cdot V \cdot \text{Accuracy}_0 \cdot \gamma)}$) connects retraining frequency to quantitative economics under its stated assumptions.
- **Deploy through graduated rollout with pretested rollback:** Canary, blue-green, and shadow deployments match risk profiles, with tiered rollback strategies that must be tested regularly through fire drills.
- **Stage the investment:** Monitoring and continuous integration/deployment are often the first investments. High-value or high-risk systems justify more rigor than low-impact internal analytics. Add feature stores when training-serving skew becomes measurable; add automated retraining as the model matures.
- **The five principles transfer across domains:** Reproducibility, separation of concerns, consistency, observable degradation, and cost-aware automation remain stable; domain constraints change their implementation and supporting infrastructure.

- **Operational maturity is staged and organizational:** Managing one model differs qualitatively from managing many. The principles scale, but complexity grows with fleet size, and shared on-call rotations and unified incentives are as critical as tooling.


The operational discipline examined in this chapter distinguishes production ML systems from development prototypes. These principles help practitioners diagnose whether degradation arises from data drift (check feature distributions), training-serving skew (compare preprocessing paths), configuration debt (audit recent changes), or feedback loop contamination (analyze temporal patterns). Teams that treat production ML as “deploy and forget” may allow models to remain wrong for months while availability dashboards stay green. As ML systems support decisions from loan approvals to medical diagnoses, this operational discipline helps organizations deploy them responsibly at scale.

A conventional system can be perfectly available and still return incorrect results. A model adds another way to reach that state: its code can stay byte-for-byte identical while its deployment distribution or target relationship changes, so a system at full uptime can be confidently, silently wrong. That is why ML operations extends software operations. The match between model and world is not a state reached once but a cost paid continuously because the data axis of D·A·M may keep moving. Reliability must include outcomes, not uptime alone, and the most dangerous state an ML system can occupy is a green dashboard resting on a drifting model.

What's Next: From reliability to responsibility

An ML system can be efficient, scalable, and reliable yet still cause harm by amplifying bias or leaking data. Even 99.9 percent uptime and sub-10 ms latency cannot establish that its decisions are responsible. Chapter 15 addresses the final constraint: aligning technical optimization with human values so deployed systems serve the people affected by them.

Responsible Engineering




- 15.1 Responsibility as Systems Engineering
- 15.2 Engineering Responsibility Gap
- 15.3 Responsible Engineering Checklist
- 15.4 Environmental and Cost Awareness
- 15.5 Data Governance and Compliance
- 15.6 Fallacies and Pitfalls
- 15.7 Summary

Purpose

Why can a system that does exactly what it was told to do still cause harm?

Operations targets low latency, high availability, and sustained predictive quality. Responsible engineering asks whether those targets describe the right behavior, whom the system serves, and which costs its specification leaves unmeasured. An ML system can satisfy latency, throughput, and aggregate accuracy requirements while reproducing historical discrimination, rewarding harmful engagement, consuming unjustified energy, or using data without adequate privacy and accountability controls. Such a system has not malfunctioned; it is efficiently optimizing an incomplete specification. The omitted constraints appear as uneven outcomes across populations, lifetime cost and emissions, decisions people cannot understand or contest, and evidence that cannot be reconstructed after harm occurs. If those consequences are not defined as requirements before deployment and monitored afterward, conventional health checks can remain green while harm accumulates. Responsible engineering treats them as system failures to be diagnosed, measured, and mitigated with the same rigor as latency or accuracy regressions. That work requires translating broad obligations into testable bounds, documented assumptions, traceable data, and response paths that remain enforceable in production. In D·A·M terms, responsibility broadens the definition of correctness: data must be examined for encoded harms and lawful use, algorithms must be bounded by defensible outcome and efficiency requirements, and machine infrastructure must monitor, document, and enforce those boundaries throughout the system lifecycle.





Learning Objectives

- Explain how optimized ML systems can amplify harm through proxies, feedback loops, and distribution shift
- Apply data-algorithm-machine diagnosis to localize responsibility failures in data, algorithm objectives, or monitoring infrastructure
- Calculate fairness metrics from confusion matrices and compare trade-offs on the fairness-accuracy Pareto frontier
- Design disaggregated evaluation, stress testing, and monitoring to expose subgroup-specific failures before deployment
- Analyze total cost, inference dominance, and carbon impact as measurable responsibility constraints
- Construct model cards, datasheets, lineage records, and audit trails for accountability
- Evaluate privacy, access-control, and compliance designs against regulatory and human-review requirements

15.1 Responsibility as Systems Engineering

¹ | **Amazon recruiting tool:** Developed starting in 2014 by Amazon's Edinburgh engineering team to rate applicants on a 1–5 scale, the system trained on approximately a decade of resumes—overwhelmingly from male applicants reflecting the tech industry's gender ratio. By 2015 the gender bias was identified; by 2017 the project was abandoned after repeated remediation attempts (Dastin 2018). The engineering cost was not the compute but the opportunity cost: a multi-year recruiting project failed because the objective encoded historical bias, making it a documented specification failure in ML tooling.

IN 2014, Amazon built an AI recruiting tool¹ that penalized resumes containing the word “women’s” (as in “women’s chess club captain”) and downgraded graduates of all-women’s colleges. The system optimized faithfully for its stated objective: identify candidates similar to those previously hired. The failure was not that the model malfunctioned, but that historical hiring patterns encoded gender bias, and the model reproduced that bias at scale.

If MLOps is the control loop for *reliability*, then **responsible engineering** is the control loop for *safety*. Where MLOps monitors system health and triggers retraining when performance degrades, responsible engineering monitors outcome quality and triggers intervention when systems cause harm. A model can optimize flawlessly for its stated objective and still cause systematic harm because the failure is not a bug in the code but a flaw in the specification. In systems engineering terms, a system can pass *verification* (it meets its stated requirements) while failing *validation* (it does not meet the user’s true needs) (National Aeronautics and Space Administration 2016).

Traditional software engineering often isolates defects within modules, although shared dependencies can still propagate failures. Machine learning systems add data-dependent coupling: data flows through shared representations, so problems in one component can affect many outputs. A biased training dataset does not produce a localized code defect; it can bias predictions across the system. Appendix A formalizes the diagnostic framework that locates where such a failure originates, decomposing it along three axes: biased data, a misaligned algorithm, or inadequate infrastructure for monitoring outcomes. This makes responsibility an architectural concern, not an afterthought.

Engineering responsibility therefore expands what “correct” means for ML systems. Correctness in the traditional sense (reliable, performant, and maintainable) remains necessary, but ML systems must also be correct in a broader sense: fair across user groups, efficient in resource consumption, and transparent in their decision processes. Expanded correctness is engineering itself, applied to failure modes that conventional metrics do not capture. A latency regression is visible in dashboards; a fairness regression is invisible until it harms real users (principle 13). Both require systematic detection, measurement, and remediation.

Diagnosing, preventing, and mitigating these failures requires following the responsibility gap through the system. Concrete cases reveal the distance between technical performance and responsible outcomes, and the mechanisms (proxy variables, feedback loops, distribution shift) through which it manifests. That gap motivates repeatable engineering processes for impact assessment, model documentation, disaggregated testing, and incident response. The resource consumption quantified throughout this book (training compute, inference energy, carbon footprint) then becomes an ethical constraint as well as a performance constraint, because efficiency optimization serves responsibility as directly as it serves speed. Data governance and compliance infrastructure

(access control, privacy protection, lineage tracking, and audit systems) make those practices enforceable at scale.

15.2 Engineering Responsibility Gap


A loan model that approves 95 percent of qualified majority-group applicants while rejecting 40 percent of equally qualified minority-group applicants can still achieve a low aggregate loss. The **responsibility gap** between this *technical correctness* and *responsible outcomes* represents a central challenge in machine learning systems engineering, one that existing testing methodologies were not designed to address. The gap manifests through concrete mechanisms: proxy variables, feedback loops, and distribution shift, each producing harm through a distinct pathway that conventional monitoring leaves invisible.

15.2.1 When optimization succeeds but systems fail

The recruiting-tool failure turns this gap into a data problem rather than a code defect. A model trained on a decade of historical hiring data optimized faithfully for the objective it was given, but those historical patterns encoded gender bias that the system reproduced in candidate ratings (Dastin 2018).

The technical mechanism behind this outcome is straightforward. The model learned token-level patterns from historical data. When most previously successful hires were men, resumes containing language associated with women’s activities or institutions appeared statistically less correlated with positive hiring decisions. The model correctly identified these patterns in the training data but learned the wrong lesson from correct pattern recognition. More generally, learned text representations can encode and amplify gender stereotypes, including in word embeddings (Bolukbasi et al. 2016).

Amazon attempted remediation by removing explicit gender indicators and gendered terms from the training process. This intervention did not establish that the remaining recommendations were unbiased because other resume features could retain sex-correlated signal.² In general, **proxy variables** can carry indirect demographic signal: ZIP codes can correlate with race because of residential segregation, names can correlate with gender or ethnicity, and healthcare utilization can correlate with socioeconomic conditions. In Amazon’s case, the reported penalties for terms such as “women’s” and for graduates of two all-women’s colleges showed how resume text could encode sex-linked patterns (Dastin 2018). Removing protected attributes from training data is therefore insufficient to establish fairness.



The Amazon failure is a data-axis failure: biased historical signal.

📖 War Story 15.1: The COMPAS recidivism algorithm audit (2016)

Context: COMPAS is a risk assessment tool used in US courtrooms to predict re-offending for bail and sentencing decisions (Angwin et al. 2016).

Mechanism: Optimizing for calibration across populations with differing baseline recidivism rates forced mathematical trade-offs between predictive parity and error-rate balance across racial groups.

Impact: Black defendants who did not re-offend were incorrectly flagged as high-risk at nearly twice the rate of White defendants (44.9 percent vs. 23.5 percent), while White recidivists were far more often mislabeled low-risk.

Response: The audit exposed why calibration alone could not settle whether the error distribution was acceptable. Any jurisdiction using such scores needs an explicit policy for which fairness criteria matter, independent validation across groups, and review of how the score enters the decision workflow.

Systems lesson: The scores were approximately calibrated by race, but they violated equalized odds because false positive and false negative rates did not match across groups. Formal fairness results show that calibration and error-rate parity conflict when base rates differ, showing why engineering responsibility requires explicitly choosing which fairness constraint matters.

² | **Proxy variable:** Removing one proxy may have little effect when other correlated features retain similar signal. Conversely, correlation alone does not prove discrimination or identify the appropriate intervention. Protected-attribute removal, subgroup evaluation, counterfactual tests, causal analysis, and review of the full decision process provide complementary evidence; no single diagnostic is a complete defense.

The right intervention would have required multiple levels of change. Separate evaluation of resume scores for male-associated vs. female-associated candidates could have revealed the disparity quantitatively. Fairness constraints or adversarial debiasing, where an auxiliary adversary tries to recover protected-attribute signal from the learned representation and the main model is penalized when that signal remains, might have reduced some encoded signal but would still require subgroup validation. Human review for borderline cases could have added a safeguard if reviewers were trained, empowered, and monitored rather than asked to rubber-stamp scores. Tracking actual hiring outcomes by gender over time would have enabled outcome monitoring beyond model metrics alone. Amazon eventually scrapped the project after determining that sufficient remediation was not feasible (Dastin 2018).

The Amazon case demonstrates how optimization objectives diverge from organizational values. The system found genuine statistical patterns in historical hiring decisions and optimized them faithfully. Those patterns, however, reflected biased historical practices rather than job-relevant qualifications.

The Amazon and COMPAS³ cases share a troubling pattern: each system achieved its stated objective while producing outcomes that conflicted with the values the system was intended to serve. Conventional engineering success can coexist with serious system failures. The pattern raises two design questions: whether the loss function is a defensible proxy for the system's true goal, and whether error rates remain acceptable across the subgroups the system affects.

Conventional tests focused only on the stated objective may not catch these problems because they represent failures of problem specification, where the *technical objective* (minimizing prediction error on historical outcomes) diverges from the *desired social objective* (making fair and accurate predictions across demographic groups). Specification failures are difficult to detect precisely because the systems continue functioning normally by conventional engineering metrics. When a system appears healthy by its monitored technical metrics, harm outside those metrics can remain invisible.

³ | **COMPAS (Correctional Offender Management Profiling for Alternative Sanctions):** In the analyzed data, COMPAS scores were approximately calibrated by race, meaning a given score corresponded to a similar observed re-offense probability across groups. Because recidivism prevalence differed between populations, calibration coexisted with disparate error rates (Chouldechova 2017; Kleinberg et al. 2017). Calibration testing alone therefore could not establish error-rate parity or determine whether the allocation of errors was acceptable.



Checkpoint 15.1: Responsible design

Responsibility is a system property, not a model property.

Failure modes

- ☐ **Alignment:** test whether the loss function is a defensible proxy for the system's true goal.
- ☐ **Disparate impact:** compare error rates across subgroups rather than relying on aggregate accuracy.

Check

- ☐ **Pre-mortem:** before deploying, ask: "If this system causes a scandal in six months, what likely went wrong?"


15.2.2 Silent failure modes

Consider a hospital sepsis model that begins recommending aggressive treatments for low-risk patients after an electronic health record (EHR) workflow change alters how vital signs are recorded. No alarm triggers: the model's confidence scores remain high, its latency stays within its service level agreement, and all system health checks pass green. The failure is silent: the input data distribution has shifted, but the monitoring pipeline has no mechanism to detect distributional drift.



Definition 15.1: Distribution shift

Distribution shift, introduced in chapter 1 and operationalized for drift detection in chapter 14, occurs when the deployment distribution differs from the distribution used to develop or evaluate the model. Standard empirical-risk minimization commonly assumes matching training and deployment distributions. Covariate shift changes $p(x)$ while holding $p(y | x)$ stable; label shift changes $p(y)$ while holding $p(x | y)$ stable; concept drift changes $p(y | x)$.



One upstream change; many silently affected.

1. **Significance:** Divergence statistics such as Jensen-Shannon divergence $\mathcal{D}_{\text{JS}}(P_t \| P_0)$ measure distributional change, not accuracy degradation. Useful alert thresholds must be calibrated empirically for each task, representation, label process, and deployment environment, then related to observed outcomes when labels are available. A $\mathcal{D}_{\text{JS}}$ value of 0.1 may be harmless for one feature space and severe for another, and input-divergence monitoring may miss concept drift.
2. **Distinction:** Distribution shift describes a change between development and deployment conditions. It does not imply that the learned mapping was correct at training time, nor does every shift reduce performance.
3. **Common pitfall:** Distribution shift is an umbrella term with several overlapping taxonomies. Monitoring only $p(x)$ can detect some covariate changes but cannot establish whether $p(y | x)$ remains stable; detecting concept drift generally requires labeled outcomes or justified proxies.

This sepsis scenario illustrates a class of failure that availability-focused monitoring is poorly equipped to handle. Some software failures are loud: a null pointer exception crashes the program, and a network timeout returns an error code. Conventional software can also return incorrect results silently. ML systems add statistical failure modes in which degraded predictions look like normal predictions. One mechanism behind this silent degradation is distribution shift.

Matching training and deployment distributions is a standard assumption in supervised-learning analysis, although adaptation methods relax it under additional assumptions. Distribution shift can also affect groups unequally, allowing aggregate metrics to mask subgroup harm.

Systems Perspective 15.1: The alignment gap

A model optimizes a proxy metric (Clicks) because the true metric (User Satisfaction) is unobservable, and the two can diverge sharply. Goodhart's Law warns that optimizing a measure can weaken its relationship to the underlying goal. In an illustrative system, $\text{Correlation}(\text{Clicks}, \text{Satisfaction})$ might begin at 0.8; once ranking maximizes clicks, it may find clickbait with high clicks but low satisfaction, and the correlation could fall to 0.2.


Conceptually, assuming normalized metrics on a common scale, equation 15.1 captures the gap:

$$\text{Gap} = \mathbb{E}[\text{Proxy}] - \mathbb{E}[\text{True}] \quad (15.1)$$

If the model increases Clicks by 20 percent but decreases Satisfaction by 5 percent, the signed alignment gap increases.

Systems insight: Engineers cannot directly optimize an unobserved goal. They need defensible proxy validation; randomized holdouts can estimate causal effects only when their outcomes measure the goal or a justified surrogate.

Distribution shift explains *why* models degrade over time (the operational detection and monitoring strategies for drift are covered in chapter 14). The failure is environmental: the world changed after the model was trained, and the model has no mechanism to notice. Retraining on fresh data can partially address this class of failure, but it cannot address a second mechanism for silent failure that operates even when the data distribution is perfectly stable. Metric misalignment occurs when the quantity the model optimizes diverges from the outcome the organization actually values. The dynamics of that divergence are made precise by Goodhart's Law: once a proxy becomes the optimization target, it stops tracking the goal it was chosen to represent.



🔍 Systems Perspective 15.2: The D-A-M taxonomy

When a system causes harm, use the D-A-M taxonomy to identify the root cause. Responsibility failures are rarely “algorithm bugs”; they are structural flaws along one of the three axes:

- **Data (information):** Does the training data reflect historical bias? (for example, Amazon’s recruiting tool learning from biased history). The failure is in the *Fuel*.
- **Algorithm (logic):** Does the objective function optimize a proxy for harm? (for example, optimizing “engagement” amplifies polarization). The failure is in the *Blueprint*.
- **Machine (physics):** Does the energy cost justify the societal benefit? (for example, training a massive model for a trivial task). The failure is in the *Engine*.

Locating the dominant failure in the taxonomy identifies the first remediation to test: better curation (Data), safer objectives (Algorithm), or more efficient infrastructure (Machine). Real failures can span axes, so the intervention must still be evaluated end to end.

The alignment gap illustrates a failure that originates in the algorithm axis: the optimization objective is misspecified, so even a model that generalizes flawlessly to new data can drive outcomes that conflict with organizational or societal goals. Distribution shift, by contrast, originates in the data axis: the inputs changed, and the learned mapping no longer reflects reality. Both failures are silent, but they demand different remediations (better objectives vs. better monitoring), and conflating the two wastes engineering effort on the wrong fix. The D-A-M taxonomy introduced in chapter 1 maps each failure to the axis it originates from (Data · Algorithm · Machine), which Appendix A defines in full.

While the D-A-M taxonomy helps diagnose *where* failures originate, engineers also need a framework for understanding *when* and *how* different failure types manifest. Table 15.1 complements that diagnostic framework by categorizing failures by detection time, spatial scope, and remediation requirements. Crashes and performance degradation trigger immediate alerts through existing infrastructure. Data quality issues, distribution shifts, and fairness violations require specialized detection mechanisms because the system continues operating normally from a technical perspective while producing increasingly problematic outputs.

Table 15.1: Illustrative ML System Failure Mode Taxonomy: Detection time, scope, and remediation depend on the application; the entries show representative cases. Silent failures such as data quality issues, distribution shift, and fairness violations may require proactive monitoring because they do not necessarily trigger traditional alerts.

Failure Type	Illustrative Detection Time	Illustrative Scope	Possible Response	Example
Crash	Immediate	Complete	Restart after correcting cause	Out of memory error
Performance Degradation	Minutes	Complete	Relieve resource contention	Latency spike from resource contention
Data Quality	Hours–days	Partial	Data correction may be needed	Corrupted inputs from upstream system
Distribution Shift	Days–weeks	Partial or all	May require adaptation	Population change due to new user segment
Fairness Violation	Weeks–months	Subpopulation	May require redesign	Bias amplification in historical patterns

The YouTube recommendation feedback loop (examined as a technical debt pattern in section 14.3.6.1) illustrates this pattern at scale (M. H. Ribeiro et al. 2020).⁴ M. H. Ribeiro et al. (2020) audited radicalization pathways on YouTube, finding migration from milder to more extreme channel categories and recommendation reachability between those categories. The broader systems lesson is that feedback loops can work exactly as designed while producing outcomes that conflict with societal values. Recommendation objectives must therefore be tested against downstream harms, not only against engagement proxies.

The News Feed case adds a twist the YouTube loop did not: a platform choosing to trade measured engagement for long-term welfare; direct evidence that engagement proxies are known to be

⁴ | **Goodhart’s law:** “When a measure becomes a target, it ceases to be a good measure” (Strathern’s generalization of Goodhart’s 1975 monetary policy observation) (Strathern 1997; Goodhart 1984). Recommendation feedback loops are the canonical ML manifestation: gradient descent optimizes watch-time proxies at a speed no human curator can match, and the system’s own outputs reshape the training distribution—users who consume extreme content generate data that reinforces extremity, decoupling the proxy from user welfare orders of magnitude faster than manual editorial processes ever could.

incomplete even by those who optimize them. A distinct failure mode operates at the population level: proxy variables that appear neutral in the aggregate can encode systematic disparities across demographic groups. The distribution shift defined in this section also manifests as population mismatch, where models trained on one population perform differently on another without obvious indicators. The same proxy mechanism that let Amazon's recruiting model reconstruct gender reappears in healthcare, where cost as a stand-in for need produced one of the most widely studied cases of algorithmic harm.

Example 15.1: News Feed proxy shifts

Scenario: Facebook shifted News Feed ranking objectives from short-term engagement proxies (clicks, watch time) to meaningful social interaction metrics (Mosseri 2018).

Diagnosis: Optimizing solely for short-term engagement proxies decoupled ranking outputs from long-term user welfare, amplifying clickbait and passive consumption.

Systems lesson: Engagement metrics are proxies for value, not value itself. Recommendation algorithms require multi-objective optimization with explicit safety constraints to prevent proxy collapse.

Silent failure modes complicate testing. Traditional software testing often verifies specified behavior through exact assertions. ML systems add behavior learned from data and performance evaluated statistically, making correctness harder to define. The opening failures share a troubling pattern: conventional technical checks did not cover the harms that emerged. Organizations reduce this gap when they convert responsibility goals into structured engineering practice.

War Story 15.2: The proxy variable trap (2019)

Context: In 2019, Ziad Obermeyer and colleagues at UC Berkeley audited a commercial Optum algorithm used by health systems to enroll patients with complex needs into high-risk care management programs, examining roughly fifty thousand patients across one large academic hospital (Obermeyer et al. 2019).

Mechanism: The model predicted “future healthcare cost” as a proxy for “future health need.” Because the US healthcare system spends less on Black patients than on White patients with the same illness level, the algorithm learned this pattern and assigned lower risk scores to Black patients.

Impact: At any given risk score, Black patients carried substantially more chronic conditions than White patients, reducing the share of Black patients enrolled in specialized care.

Response: The study showed that replacing cost with a direct measure of health need would raise the share of Black patients identified for additional care from 17.7 percent to 46.5 percent at the analyzed threshold.

Systems lesson: Optimizing for a proxy inherits the biases of the system that generated the proxy. The proxy-target relationship must be audited across every demographic subgroup the system serves.

15.2.3 When responsible engineering succeeds

Each documented success shares the same structural move: a vague responsibility goal becomes an engineering constraint that can be specified, tested, communicated, and, when necessary, used to stop deployment. Following the findings of Gender Shades, a 2018 audit that exposed severe error-rate disparities in commercial gender classification (Buolamwini and Gebru 2018), Microsoft invested in improving gender-classification performance across demographic groups. Targeted data collection, model changes, and systematic disaggregated evaluation gave the team an explicit error-rate target, and Microsoft reported large error-rate reductions for darker-skinned subjects, bringing audited error rates below 2 percent (Raji and Buolamwini 2019). The company published these improvements transparently, turning external audit results into measurable engineering targets.

Twitter's automatic image cropping system shows the same discipline under a different constraint. In 2020, users raised concerns about racial and gender bias in which faces appeared in preview thumbnails. Twitter audited the system, reported measured disparities and limitations, and shifted toward uncropped previews that gave users more control (Yee et al. 2021). In that case, responsible engineering changed the product design rather than relying only on a better ranking threshold.

Differential privacy makes the pattern formal: a privacy requirement becomes a mathematical guarantee rather than a policy aspiration (Dwork 2008).⁵ Systems that use differential privacy must calibrate noise to balance utility against privacy, track privacy budget across repeated analyses, and document the chosen privacy parameters.

A common pattern unites these cases: responsibility creates value only when technical interventions (improved data, better evaluation, architectural changes, formal guarantees, or user controls) combine with organizational commitments to transparency, long-term investment, and the willingness to remove features that cannot be made safe. Each success rested on systematic testing and evaluation practices, yet the nature of responsible testing differs fundamentally from traditional software verification.

15.2.4 The testing challenge

Traditional software testing verifies that systems behave correctly because correctness has clear definitions. The function should return the sum of its inputs, the database should maintain referential integrity. These properties can be expressed as testable assertions.

Responsible ML properties resist simple formalization. **Fairness** has multiple mathematical definitions that can conflict under conditions such as unequal base rates and imperfect prediction (Chouldechova 2017; Kleinberg et al. 2017). What counts as fair depends on context, values, and trade-offs that technical systems cannot resolve alone. Individual fairness requires that similar individuals receive similar treatment, while group fairness requires equitable outcomes across demographic categories. These criteria can conflict, and choosing between them requires value judgments beyond the scope of optimization.

Fairness constraints and predictive performance can trade off, but the relationship is application-specific. A Pareto frontier represents configurations for which one plotted objective cannot improve without degrading another. Figure 15.1 visualizes one hypothetical fairness-accuracy frontier. Its three points illustrate possible policy choices; they do not imply that unconstrained training always maximizes accuracy at high disparity, that zero disparity must reduce accuracy, or that every application has a sweet spot.




Figure 15.1: An Illustrative Fairness-Accuracy Pareto Frontier: A hypothetical relationship between model accuracy and demographic disparity. The x-axis is inverted so lower disparity lies to the right. Points A, B, and C illustrate three possible operating choices along this constructed frontier; the curve is not an empirical or universal law.

The frontier tells engineers what trade-off they may need to choose, but it cannot be plotted until subgroup performance is measured. Responsible properties become testable when engineers work

⁵ | **Differential privacy:** Introduced by Dwork et al. (2006), a randomized mechanism $\mathcal{M}$ satisfies (ϵ, δ) -differential privacy if for all neighboring datasets D, D' differing by one record and output sets $\mathcal{S}$, $\mathbb{P}[\mathcal{M}(D) \in \mathcal{S}] \leq e^\epsilon \cdot \mathbb{P}[\mathcal{M}(D') \in \mathcal{S}] + \delta$. Here ϵ bounds privacy loss and δ permits a small probability of exceeding that bound; their acceptable values depend on the threat model and policy. The systems trade-off is utility rather than mere implementation complexity: stronger privacy often requires more noise or tighter sampling, and privacy loss composes across repeated queries or training steps. Engineers must therefore track cumulative privacy loss with a valid accountant and report the assumptions behind the chosen budget.

with stakeholders to define criteria appropriate for specific applications. The Gender Shades project⁶ demonstrated how **disaggregated evaluation** across demographic categories reveals disparities invisible in aggregate metrics (Buolamwini and Gebru 2018), exposing the subgroup failures that responsibility monitoring must catch before deployment. Table 15.2 shows the dramatic error-rate differences that commercial gender-classification systems produced across demographic groups. Concretely, a 10,000-sample test set that suffices for the majority group provides only 100 samples for a minority subgroup representing 1 percent of the population—effectively requiring 100× more data than the majority group for high-confidence validation.

Table 15.2: Gender Shades Face++ Error Rates: The dark-female error rate is 43.1× the light-male rate. Source: (Buolamwini and Gebru 2018).

Demographic Group	Error Rate (%)	Relative to Light-Skinned Males
Light-skinned males	0.8%	Baseline (1.0×)
Light-skinned females	9.8%	12.2×
Dark-skinned males	0.7%	0.9×
Dark-skinned females	34.5%	43.1×

As table 15.2 quantifies, disaggregated evaluation revealed what an aggregate score concealed. Face++ reported error rates of 0.8 percent for light-skinned males and 34.5 percent for dark-skinned females (accuracies of 99.2 percent and 65.5 percent). The aggregate metric gave no indication that the dark-female error rate was 43.1× the light-male rate.

No universal threshold defines acceptable disparity, but teams should establish explicit, justified bounds before deployment. A team might adopt an error-rate ratio below 1.25× or a false-positive-rate difference under 5 percentage points as an application-specific release policy. In US employment law, the disparate impact doctrine⁷ shapes regulatory oversight, while the four-fifths rule⁸ provides a separate, non-dispositive selection-rate diagnostic. The key engineering discipline is defining criteria appropriate to the application and governing law rather than generalizing one domain's threshold.

Despite the inherent challenges, several concrete testing approaches can surface responsibility issues before deployment:

- **Slice-based evaluation:** Partitions test data into meaningful subgroups and reports metrics separately for each slice, asking whether aggregate performance hides subgroup failure. A model may achieve 95 percent accuracy overall but only 78 percent accuracy on low-income applicants or users from rural areas, a disparity invisible in aggregate reporting.
- **Invariance testing:** Checks whether the model changes behavior for the wrong reasons. Replacing “John” with “Jamal” in a loan application should not change approval likelihood if the feature is not legitimate for the decision. Behavioral testing frameworks such as CheckList apply this idea by organizing tests around model capabilities and invariance-style expectations rather than accuracy alone (M. T. Ribeiro et al. 2020).
- **Boundary and stress testing:** Probes regions where ordinary validation sets are least informative. Boundary testing evaluates model behavior at the edges of input distributions (unusual ages, extreme values, rare categories) where training data may be sparse and predictions unreliable. Stress testing extends boundary testing to adversarial conditions: corrupted inputs, distribution shift, adversarial examples, and edge cases designed to probe failure modes systematically. Stakeholder red-teaming adds evidence from domain experts and affected community members, surfacing failure modes no automated test can discover because they require lived experience to imagine.

Responsible testing strategies complement traditional software testing rather than replacing it. Each demands joint judgment to select, configure, and interpret. Legal specialists cannot alone specify which demographic slices matter for a healthcare algorithm, and product managers cannot alone determine appropriate invariance tests for a loan model. Domain experts, affected stakeholders, policy specialists, and engineers must define the relevant harms together; engineers then encode

⁶ | **Gender Shades:** A 2018 study by Joy Buolamwini (MIT Media Lab) and Timnit Gebru (Microsoft Research) that audited commercial gender-classification systems from Microsoft, IBM, and Face++ using the Fitzpatrick skin type scale—originally a dermatological classification developed by Thomas Fitzpatrick in 1975 for UV sensitivity and later validated for clinical use (Fitzpatrick 1988), repurposed here as a demographic benchmark for algorithmic auditing. The study demonstrated the value of disaggregated evaluation; in the Face++ results reproduced in table 15.2, the dark-female error rate is 43.1× the light-male rate. Microsoft’s later reported reductions show how public audit results can motivate measurable remediation (Raji and Buolamwini 2019).

⁷ | **Disparate impact:** In *Griggs v. Duke Power Co.* (1971), the US Supreme Court held under Title VII that an employment practice may be unlawful because of its operation even without discriminatory intent (Supreme Court of the United States 1971). Disparate impact is a legal claim with statute- and context-specific elements, not a synonym for any statistical disparity. Model outcomes and proxies can supply relevant evidence, but metrics alone do not establish liability.

⁸ | **Four-fifths rule:** The 1978 Uniform Guidelines on Employee Selection Procedures use the four-fifths rule as a practical employment-selection diagnostic (Equal Employment Opportunity Commission et al. 1978). A selection rate below 80 percent of the rate for the group with the highest rate is generally regarded as evidence of adverse impact, but the Guidelines also recognize that smaller differences may matter and that the ratio is not dispositive.

those requirements as measurable tests and production controls. Responsibility therefore needs clear ownership within engineering as well as independent oversight outside it.

15.2.5 Engineering leadership on responsibility

By the time Amazon abandoned the recruiting tool, attempted remediation had not provided confidence that it would avoid discriminatory recommendations (Dastin 2018). Earlier design choices can narrow the available fixes. Responsible AI engineering cannot be delegated exclusively to ethics boards or legal departments. These groups provide essential oversight but lack the technical access required to identify problems early in the development process.

Legal or ethics review can identify a problem near deployment, but it cannot recover design options the system has already foreclosed. If the team trained the model without fairness constraints, chose an architecture that cannot support interpretability requirements, or built a data pipeline without the demographic attributes needed for monitoring, review can only accept, reject, or demand expensive redesign. Engineers therefore occupy a critical position in the ML development lifecycle because their choices define the solution space for all subsequent interventions: architecture determines which fairness constraints can apply, the optimization objective determines which patterns the system learns, and the data pipeline determines whether disaggregated evaluation is possible.



Definition 15.2: Responsible AI engineering

Responsible AI engineering is the engineering discipline of designing, deploying, and maintaining systems with probabilistic outputs by operationalizing societal and regulatory requirements as testable constraints on the D·A·M axes: permissible data contents, provenance, and composition; allowable model behavior and robustness properties; and infrastructure bounds such as latency, energy, compute budget, carbon emissions, and audit-log retention.

1. **Significance:** Each D·A·M axis acquires concrete governance constraints: the data axis is bounded by privacy regulations such as the General Data Protection Regulation (GDPR), which limits which records, fields, and features can be collected; the algorithm axis is bounded by fairness and robustness metrics (for example, demographic parity within $\varepsilon = 5\%$ across protected groups, meaning positive prediction rates must not differ by more than 5 percentage points, or accuracy degradation less than 2 percent under adversarial perturbation $\|\delta\|_\infty \leq 0.01$, a worst-case input change bounded to 0.01 per normalized feature under the ℓ_∞ norm); and the machine axis is bounded by resource and infrastructure budgets such as latency, energy per inference, carbon emissions, and audit-log retention. Violating these bounds is a system failure, not a research shortcoming.
2. **Distinction:** Unlike AI ethics (which articulates aspirational values), responsible AI engineering translates those values into measurable, testable invariants that can be verified through automated testing and continuous monitoring, using the same lifecycle practices that enforce latency SLOs.
3. **Common pitfall:** A frequent misconception is that responsibility is “added” at the end of development. Constraints on what data may be collected affect what can be learned and audited, while infrastructure choices affect what evidence can be retained. Late-stage remediation may remain possible, but it is often narrower, slower, and more expensive.

An engineering-centered approach does not diminish the importance of diverse perspectives in identifying potential harms. Product managers, user researchers, affected communities, and policy experts contribute essential knowledge about how systems fail socially despite technical success. Engineers translate these concerns into measurable requirements and testable properties that can be verified throughout the development lifecycle. Effective responsibility requires engineers who both listen to stakeholder concerns and possess the technical capability to implement appropriate safeguards.

Engineering teams do not operate in isolation. As figure 15.2 makes clear, engineering practices are nested within broader organizational, industry, and regulatory governance structures, each layer

imposing constraints on the ones inside it. Technical excellence at the innermost layer enables, but does not replace, compliance with requirements flowing inward from external governance.




Figure 15.2: Responsible AI Governance Layers: Four nested ovals place team-level safeguards at the center, surrounded by organizational safety culture, industry certification and review, and government regulation. The nesting distinguishes implementation at the center from oversight outside it; the annotation states that requirements flow inward and technical practices enable compliance. Adapted from Shneiderman (2022).

Those governance layers define who owns responsibility, but they do not yet account for the costs that ordinary performance metrics omit.

Beyond ethical imperatives, responsible engineering can deliver business value through three reinforcing mechanisms. The most immediate is risk mitigation: ML system failures create legal and financial exposure that systematic responsibility practices can reduce. Amazon abandoned an internal recruiting model after discovering sex-linked disparities in its recommendations. Organizations implementing disaggregated evaluation, documentation, and monitoring can reduce the probability of costly failures and preserve evidence of their engineering process if problems emerge.

A second mechanism is regulatory compliance, driven by legal requirements that vary by jurisdiction and application risk. The EU AI Act, for example, classifies high-risk AI applications and mandates technical requirements including risk assessment, data governance, transparency, and human oversight. Organizations that build responsibility into engineering practice can demonstrate compliance through existing documentation and monitoring rather than expensive retrofitting; the engineering lesson is that proactive controls are usually cheaper than reconstructing evidence after deployment.

Systems Perspective 15.3: The full cost of the iron law

The iron law of ML systems (principle 3) established in section 1.7 holds that system performance depends on the interaction between data, compute, and system overhead. Earlier chapters optimized each term by compressing models (chapter 10), accelerating hardware (chapter 11), and automating operations (chapter 14). Yet an optimization can create or shift costs that its benchmark omits.

A model quantized for edge deployment consumes less energy, but also produces outputs that may differ across demographic groups. A recommendation system optimized for engagement maximizes a business metric, but may amplify harmful content. Responsible engineering extends our accounting to include these broader impacts: the carbon cost of computation, the

fairness cost of optimization choices, and the societal cost of deployment at scale. The iron law governs *how fast* our systems run; responsible engineering governs *how well* they serve.

Competitive differentiation completes the business case. Trust can drive enterprise purchasing decisions for ML-powered services, and organizations that can demonstrate systematic responsibility practices through model cards, audit trails, and published evaluation results may qualify for deployments that competitors cannot. Apple’s privacy positioning, Microsoft’s responsible AI principles, and Anthropic’s safety research illustrate responsibility as a strategic investment rather than a purely defensive cost.

The quantization techniques from chapter 10 can reduce inference energy when the deployed runtime and hardware exploit the representation and measured workload energy falls. The monitoring infrastructure from chapter 14 can support disaggregated fairness evaluation when the system lawfully collects the necessary outcome and subgroup data. Responsible engineering synthesizes these capabilities into disciplined practice through structured frameworks that translate principles into processes.

Systematic processes applied early could have reduced the likelihood or severity of the failures examined in section 15.1. Checklists, documentation standards, testing protocols, and monitoring infrastructure translate responsibility principles into repeatable engineering workflows, but they do not guarantee prevention.

15.3 Responsible Engineering Checklist

A structured predeployment review could have surfaced the sex-linked signals learned by Amazon’s recruiting tool, while disaggregated testing exposed COMPAS’s error-rate disparity. Both failures shared a common cause: responsibility was treated as a separate review stage rather than integrated into the development workflow. A responsible engineering checklist embeds assessment wherever engineering decisions create durable risk: before deployment, in documentation, during population-specific evaluation, at explanation and compliance boundaries, and after launch through monitoring. The stages build on one another: assessment identifies what to measure, documentation preserves the assumptions, fairness evaluation checks whether performance holds across groups, explainability and compliance translate decisions into obligations, and monitoring connects detected violations to intervention.

15.3.1 Predeployment assessment

Before a loan approval model reaches production, a team must determine the provenance of the training data, identify who is represented and who is missing, anticipate failure modes, and define recourse for affected users. Table 15.3 structures this evaluation into five phases, distinguishing critical-path blockers from high-priority items that can proceed with documented risk acceptance.

Table 15.3: Predeployment Assessment Framework: Critical Path items block deployment until addressed. High Priority items should be completed before or shortly after launch. Systematic coverage of responsibility concerns throughout the ML lifecycle reduces the chance that risks are overlooked.

Phase	Priority	Key Questions	Documentation Required
Data	Critical Path	Where did this data come from? Who is represented? Who is missing? What historical biases might be encoded?	Data provenance records, demographic composition analysis, collection methodology documentation
Training	High Priority	What are we optimizing for? What might we be implicitly penalizing? How do architecture choices affect outcomes?	Objective function specification, regularization choices, hyperparameter selection rationale
Evaluation	Critical Path	Does performance hold across different user groups? What edge cases exist? How were test sets constructed?	Disaggregated metrics by demographic group, edge case testing results, test set composition analysis
Deployment	Critical Path	Who will this system affect? What happens when it fails? What recourse do affected users have?	Impact assessment, stakeholder identification, rollback procedures, user notification protocols

Phase	Priority	Key Questions	Documentation Required
Monitoring	High Priority	How will we detect problems? Who reviews system behavior? What triggers intervention?	Monitoring dashboard specifications, alert thresholds, review schedules, escalation procedures

Critical Path items are deployment blockers: the system must not go to production until these questions are answered. High Priority items should be addressed but may proceed with documented risk acceptance and a remediation timeline. The distinction enables teams to ship responsibly without requiring perfection on every dimension before initial deployment.


The Evaluation row in table 15.3 raises the critical concern of whether performance holds across different user groups. Answering this question requires statistically valid test sets for each group, which can create surprisingly stringent data requirements when representation is uneven.

 Napkin Math 15.1: The statistics of representation

Problem: An engineering team needs to verify that a Face ID model works for a minority group representing 1 percent of the user base. A worst-case binomial margin of error near 1 percentage point at 95 percent confidence requires roughly 10,000 images for this group.

Random sampling: To get 10,000 images of a 1 percent group via random sampling, the team must collect and label: $D_{eval,total} = 10,000 \text{ images} / 0.01 = 1,000,000 \text{ images}$

Stratified sampling: Specifically targeting this group (for example, via active learning or community outreach) requires only 10,000 images. **Systems insight:** Relying on “natural distribution” data for fairness is prohibitively expensive under random sampling. Validating the minority group effectively requires 100× more data than the majority group. Fairness requires intentional data engineering, not just more data.



Random sampling barely reaches small subgroups.

Intentional data engineering addresses *what* the model sees during evaluation, but even a perfectly representative dataset cannot prevent harm at deployment if the system lacks adequate human oversight. The representation cost derived in napkin math 15.1 is a predeployment gate; the question that follows is what happens once the model is live and making decisions that affect people.

 War Story 15.3: The automation paradox (2018)

Context: Uber’s Advanced Technologies Group (ATG) was testing self-driving cars in Arizona, relying on a human safety driver to intervene if the autonomous perception system failed (National Transportation Safety Board 2019).

Mechanism: The perception system detected a pedestrian crossing the road but repeatedly toggled its classification between vehicle, bicycle, and unknown, resetting its trajectory prediction while automatic emergency braking was intentionally disabled.

Impact: The safety driver was visually distracted by a personal phone and failed to take control, resulting in a fatal collision.

Fix: Uber retrofitted the fleet with driver-monitoring cameras, re-enabled automatic emergency braking, and established dual-operator staffing during autonomous road testing.

Systems lesson: Adding a human backup creates a new system with its own failure modes. High automation reliability can encourage complacency and reduce vigilance, so an effective fallback requires workload design, attention monitoring, and a safety case for the combined human-machine system.



Without effective controls, reliable automation can erode vigilance.

For high-stakes applications, the deployment phase should specify where human oversight is required. Human-in-the-loop (HITL) systems route uncertain, high-consequence, or flagged decisions to human reviewers rather than acting autonomously. Effective HITL design must specify four requirements: the review scope (which decisions require human review), the confidence thresholds that trigger escalation, the training reviewers receive, and the mechanisms for monitoring reviewer

performance. HITL is not a catch-all solution: human reviewers can rubber-stamp automated decisions, introduce their own biases, or become overwhelmed by alert volume. Effective HITL design requires calibrating the human-machine boundary to the specific application risks and reviewer capabilities.

The predeployment assessment framework parallels aviation preflight checklists, which standardize critical checks under time pressure. Production ML deployments require equivalent discipline and rigorous verification. A checklist prompts teams to ask the right questions but does not prove that a system is safe; documentation standards preserve the answers and allow them to travel with the model.

15.3.2 Model documentation standards

Consider inheriting a production model from a departed colleague: the model achieves 94 percent accuracy on its test set, but three pieces of information are missing. The identity of that test set, the data the model was trained on, and the populations it was validated against are unknown. Without those answers, deploying or updating the model is a gamble. **Model cards** solve this problem by providing a standardized documentation format for ML models⁹ (Mitchell et al. 2019). Originally developed at Google, model cards function as “nutrition labels” that capture information essential for responsible deployment and travel with the model throughout its lifecycle.

A complete model card covers seven concerns that together enable responsible deployment. It begins with technical details (architecture, training procedures, hyperparameters) that enable reproducibility and auditing. Crucially, it specifies *intended use* alongside explicit exclusions, preventing the scope creep where models designed for photo organization get repurposed for security screening. The card then documents which *factors* (demographic groups, environmental conditions, instrumentation differences) might affect performance, guiding both evaluation strategy and monitoring protocols.

The remaining sections close the gap between what a model *can* do and what it *should* do. Performance metrics must include disaggregated results across the factors identified in section 15.2.4, because aggregate accuracy alone conceals the disparities this chapter has documented. Training and evaluation data documentation enables assessment of potential encoded biases and provides essential context for interpreting results. Ethical considerations make implicit trade-offs explicit by documenting known limitations, potential harms, and mitigations implemented, while caveats and recommendations provide guidance on appropriate use and known failure modes.

A concrete MobileNetV2 model card makes these abstract categories operational: table 15.4 shows how each section addresses specific deployment concerns for edge deployment.

Table 15.4: Illustrative Model Card: MobileNetV2 for Edge Deployment: Abstract model card categories translate to practical documentation that guides responsible deployment decisions. The numerical entries are scenario assumptions, not reported benchmark results.

Section	Content
Model Details	MobileNetV2 architecture with 3.5M parameters, trained on ImageNet using depthwise separable convolutions. INT8 quantized for edge deployment.
Intended Use	Real-time image classification on mobile devices with less than 50 ms latency requirement. Suitable for consumer applications including photo organization and accessibility features.
Factors	Performance varies with image quality (blur, lighting), object size in frame, and categories outside ImageNet distribution.
Metrics	71.8% top-1 accuracy on ImageNet validation (full precision: 72.0%). Accuracy varies by category: 85% on common objects, 45% on fine-grained distinctions.
Ethical Considerations	Training data reflects ImageNet biases in geographic and demographic representation. Not validated for high-stakes applications (medical diagnosis, security screening). Performance may degrade on images from underrepresented regions.

Datasheets for datasets provide analogous documentation for training data (Gebru et al. 2021). These documents capture data provenance, collection methodology, demographic composition, and known limitations that affect downstream model behavior. Documentation establishes what a model is designed to do; testing verifies whether it performs equitably across the populations it serves.

⁹ | **Model cards:** The primary failure mode model cards address is scope creep: gradual expansion from “it worked for case A” to “try it for case B” without revalidating intended use; in practice, cards are often written after deployment decisions are made, documenting observed behavior rather than constraining it. The companion “Datasheets for Datasets” (Gebru et al. 2021) applies the same principle to training data. Without both, the card becomes a historical record rather than a guard rail.

15.3.3 Testing across populations

The disaggregated evaluation that exposed the Gender Shades disparities (section 15.2.4) now becomes an operational release-gate task: selecting the slices, metrics, and thresholds that determine whether deployment is allowed. Aggregate performance metrics mask disparities across user populations, the **flaw of averages** (Savage 2009); responsible testing requires disaggregated evaluation that examines performance for each relevant subgroup.

Systems Perspective 15.4: The flaw of averages

In systems engineering, we rarely design for the “average” case; we design for the tail cases and boundary conditions. A bridge that is “safe on average” but collapses under a heavy truck is a failure. Similarly, an ML system that is “accurate on average” but fails for a specific ethnic or gender group is an engineering failure. The same principle that drives us to measure tail latency (p99) for system reliability applies to fairness: we must use disaggregated evaluation to measure system fairness. Looking only at aggregate accuracy blinds the analysis to systemic failures occurring in the margins. Responsible engineering requires making these “tails” visible through granular, population-specific measurement.

The flaw of averages gives the testing principle: aggregate metrics are not enough. The next question is which tails to expose. That answer depends on the workload archetype, because each archetype creates different opportunities for bias to enter and different metrics for detecting it.

Table 15.5 turns the flaw of averages into an engineering workflow: a vision model fails differently than a recommendation system, so the fairness metrics must match the failure mode and the subgroup slices must come from the application context. For healthcare applications, demographic factors like race, age, and gender are essential. For content moderation, language and cultural context matter. For financial services, protected categories under fair lending laws require specific attention.

Testing infrastructure should support stratified evaluation where performance metrics are computed separately for each relevant subgroup, enabling comparison of error rates and error types across populations. Intersectional analysis considers combinations of attributes because harms may concentrate at intersections not visible in single-factor analysis. Confidence intervals provide uncertainty quantification for subgroup metrics when small subgroup sizes may yield unreliable estimates. Temporal monitoring tracks subgroup performance over time, detecting drift that affects some populations before others.

Tool choice matters only after the team has named the fairness metric, subgroup slices, and alert thresholds. Open-source libraries such as Fairlearn (Bird et al. 2020), AI Fairness 360 (Bellamy et al. 2019), and Google’s What-If Tool (Wexler et al. 2020) lower the implementation cost of disaggregated and intersectional evaluation, but a library can only compute the metric the engineer asks for. It cannot decide which subgroup definition matters, which disparity threshold should page the team, or which fairness constraint is appropriate for the deployment context.

A release gate must distinguish absence of evidence from evidence of parity. If a slice lacks enough labeled examples to bound the relevant error rate, report insufficient data—not a pass—and collect targeted data or restrict intended use. Population testing therefore produces three operational outcomes: pass, fail, or insufficient evidence.

Lighthouse 15.1: Fairness concerns by archetype

Fairness risks vary across the workload archetypes in chapter 2. Table 15.5 maps each to a primary risk and metric.

Architecture alone does not determine fairness; each row identifies a deployment-specific risk to test first, not a universal property of the model family.

Table 15.5: Fairness Risk by ML Archetype: Fairness risks vary by archetype’s data source and deployment context.

Archetype	Primary Fairness Risk	Key Evaluation Metric	Real-World Example
ResNet-50 (Compute Beast)	Training data bias (underrepresentation of deployment-relevant groups)	Disaggregated accuracy by demographic group	Evaluate the deployed classifier on application-specific demographic slices; Gender Shades values do not measure ResNet-50
GPT-2 (Bandwidth Hog)	Corpus bias (overrepresentation of majority viewpoints in web text)	Toxicity rate by demographic prompt context; stereotype score	LLMs produce more toxic completions for prompts mentioning minority groups
DLRM (Sparse Scatter)	Feedback-loop amplification (popular items get more data)	Share of recommendation impressions by item category and supplier or creator group	Filter bubbles: the system recommends similar content to similar users, reducing discovery of niche creators
DS-CNN (Tiny Constraint)	Deployment-context mismatch (trained on clean audio, deployed in noisy real-world environments)	False positive rate by acoustic environment and speaker accent	Voice assistants perform worse on accented speech; wake-word triggers on TV audio in some languages

Systems insight: Fairness evaluation must match each archetype’s failure mode. Vision models require stratified demographic accuracy; large language models (LLMs) need toxicity and stereotype probes; recommenders need exposure audits; TinyML needs acoustic-environment tests. The lighthouse keyword spotting (KWS) system introduced in chapter 2 as the Tiny Constraint lighthouse faces exactly this challenge for its DS-CNN, a depthwise-separable convolutional neural network (CNN): trained on clean studio audio, it must perform equitably across accents, background noise levels, and speaker demographics in production homes (a governance challenge addressed in section 15.5).

15.3.3.1 Worked example: Fairness analysis in loan approval

A loan approval model reports 85 percent accuracy on the majority group and 82.5 percent overall accuracy across the evaluated applicants—numbers that may satisfy a coarse aggregate dashboard. Table 15.6 and table 15.7 reveal what the aggregate conceals: loan approval outcomes for the same model evaluated separately on two demographic groups.

Table 15.6: Confusion Matrix for Group A (Majority): In this constructed scenario, loan approval outcomes for 10,000 applicants from the majority demographic group. The 90 percent true positive rate (4,500 approved of 5,000 qualified) and 20 percent false positive rate establish the baseline for fairness comparison.

	Approved (pred)	Rejected (pred)
Repaid (actual)	4,500 (TP)	500 (FN)
Defaulted (actual)	1,000 (FP)	4,000 (TN)

Table 15.6 establishes the majority-group baseline; after the metric definitions, table 15.7 presents the minority-group comparison.¹⁰

From the same two confusion matrices, we can test three different definitions of equal treatment:

- **Demographic parity:** Requires equal positive selection rates across groups regardless of true qualification ($P(\hat{Y} = 1 \mid A = a) = P(\hat{Y} = 1 \mid A = b)$). Group A receives approval at a rate of $(4,500 + 1,000)/10,000 = 55\%$, while Group B receives approval at $(600 + 200)/2,000 = 40\%$. The 15 percentage-point disparity indicates unequal treatment in approval decisions.
- **Equal opportunity:** Requires equal true positive rates among qualified applicants ($P(\hat{Y} = 1 \mid Y = 1, A = a) = P(\hat{Y} = 1 \mid Y = 1, A = b)$). Group A achieves a TPR of $4,500/(4,500 + 500) = 90\%$, meaning 90 percent of applicants who would repay receive approval. Group B achieves only $600/(600 + 400) = 60\%$. This 30 percentage-point disparity means qualified applicants from Group B face substantially higher rejection rates than equally qualified applicants from Group A.

¹⁰ | **Fairness metric incompatibility:** The measured disparities in this worked example show how one set of confusion matrices can violate demographic parity, equal opportunity, and equalized odds. A separate impossibility result shows that, with unequal group base rates and imperfect prediction, several desirable criteria—including calibration-style predictive parity and error-rate balance—cannot generally be satisfied together (Chouldechova 2017). In those settings, optimizing one criterion can degrade another. A system designer must therefore make the trade-off explicit rather than assuming all guarantees can be achieved at once.

- **Equalized odds:** Requires both equal true positive rates and equal false positive rates ($P(\hat{Y} = 1 \mid Y = y, A = a) = P(\hat{Y} = 1 \mid Y = y, A = b)$ for $y \in \{0, 1\}$).¹¹ Group A shows an FPR of $1,000 / (1,000 + 4,000) = 20\%$, and Group B shows $200 / (200 + 800) = 20\%$. While false positive rates are equal, the true positive rate disparity means equalized odds is violated. More generally, unequal group base rates and imperfect prediction make several fairness criteria mutually incompatible, so the applicable guarantee must be chosen and justified rather than inferred from one aggregate metric.

Table 15.7: Confusion Matrix for Group B (Minority): In the same constructed scenario, loan approval outcomes for 2,000 applicants from the minority demographic group. The 60 percent true positive rate (600 approved of 1,000 qualified) reveals a 30 percentage-point disparity compared with Group A; the confusion matrices alone do not identify its cause.

	Approved (pred)	Rejected (pred)
Repaid (actual)	600 (TP)	400 (FN)
Defaulted (actual)	200 (FP)	800 (TN)

The metrics disagree because they encode different policy choices about which error rates matter most.

The metrics show that the model rejects qualified applicants from Group B at a much higher rate (40 percent false negative rate vs. 10 percent) while maintaining similar false positive rates. These confusion matrices establish a disparity, not whether it arose from thresholds, labels, features, sampling, historical discrimination, or another cause.

Where collection and use of protected attributes are lawful and justified, production systems can automate these calculations and trigger review when disparities exceed predefined thresholds. Listing 15.1 makes the failure-handling contract explicit: undefined group rates are reported as insufficient data rather than compared against a threshold.

Listing 15.1: Automated Fairness Monitoring: The core pattern computes per-group metrics from confusion matrices and alerts when disparities exceed application-specific thresholds. Production use requires sufficient data and lawful, justified handling of group attributes.

```
def compute_fairness_metrics(confusion_matrix):
    tp, fp, tn, fn = (
        confusion_matrix[k] for k in ["TP", "FP", "TN", "FN"]
    )
    total = tp + fp + tn + fn
    return {
        # Demographic parity
        "approval_rate": (tp + fp) / total if total else None,
        # Equal opportunity
        "tpr": tp / (tp + fn) if (tp + fn) else None,
        # Equalized odds (with TPR)
        "fpr": fp / (fp + tn) if (fp + tn) else None,
    }

# Compare groups and flag disparities exceeding threshold
for metric in ["approval_rate", "tpr", "fpr"]:
    if metrics_a[metric] is None or metrics_b[metric] is None:
        report_insufficient_data(metric)
        continue
    disparity = abs(metrics_a[metric] - metrics_b[metric])
    # e.g., 0.05 for high-stakes applications
    if disparity > FAIRNESS_THRESHOLD:
        trigger_alert(metric, disparity)
```

Automated monitoring achieves what manual auditing cannot at scale: continuous tracking of fairness metrics with immediate alerting when disparities emerge. The 30 percentage-point TPR

¹¹ | **Equalized odds:** Formalized by Hardt et al. (2016), requiring that both TPR and FPR be equal across protected groups; the weaker “equal opportunity” relaxes this to TPR alone. Equalized-odds post-processing may require randomized, group-dependent decisions rather than deterministic thresholds alone. It also requires protected-group information at decision time and must be evaluated under applicable law; for example, US credit regulation generally restricts using protected bases to assess creditworthiness.

disparity far exceeds this scenario’s 5-percentage-point release threshold, indicating the model requires fairness intervention before deployment. Table 15.8 reveals the pattern across the computed metrics and disparities.

Table 15.8: Fairness Metrics Summary: In this constructed loan scenario, approval rate and true positive rate differ across groups, while false positive rates match; the table establishes disparity but not its cause.

Metric	Group A	Group B	Disparity
Approval Rate	55%	40%	15 pp
True Positive Rate	90%	60%	30 pp
False Positive Rate	20%	20%	0 pp

To understand how aggregate metrics can hide disparities, look closely at figure 15.3. In this constructed example, each shared threshold applied to different score distributions produces different subgroup outcomes (Barocas and Selbst 2016). Aggregate performance alone does not determine whether those outcomes are acceptable for either group.

Several mitigation approaches exist, each with distinct trade-offs:

- **Threshold adjustment:** Lowers the approval threshold for Group B to equalize TPR but may increase false positives for that group.
- **Reweighting:** Increases the weight of Group B samples during training to give the model stronger signal about this population but may reduce overall accuracy.¹²
- **Adversarial debiasing:** Trains with an adversary that discourages the representation from encoding group membership but adds training complexity.¹³

¹² | **Reweighting:** A pre-processing technique rooted in importance sampling from statistics: samples from an underrepresented group receive higher loss weights during training, amplifying their influence on gradient updates without removing any data. Kamiran and Calders (2012) showed that appropriately chosen weights can reduce disparate impact from training data. The systems trade-off is application-specific: reweighting shifts the loss landscape and may reduce performance on other slices, so the cost must be evaluated against the Pareto frontier for the application.

¹³ | **Adversarial debiasing:** The key differentiating property is representation pressure: the adversary discourages the primary model from encoding protected-attribute information, which can reduce protected-attribute leakage and help satisfy selected fairness criteria under the evaluated distribution (B. H. Zhang et al. 2018); it does not provide a general fairness guarantee under arbitrary deployment shift; guarantees depend on assumptions about invariance, labels, causal structure, and the type of shift. Postprocessing methods such as threshold adjustment may also be appropriate under different assumptions but must be reevaluated when deployment demographics or label processes change. The cost is additional training, hyperparameter tuning, and slice-level validation rather than a universal percentage overhead.




Figure 15.3: Threshold Effects on Subgroup Outcomes: Two candidate classification thresholds intersect different score distributions for Subgroups A and B. Circles mark positive outcomes (loan repayment), plus markers mark negative outcomes (default), and color distinguishes subgroup; each boundary approves the markers to its left. Each dashed line carries the overall accuracy that boundary achieves, not a score cutoff. Moving to the higher-accuracy boundary lifts overall accuracy from 75 percent to 81.25 percent, but the whole gain falls in Subgroup A, which becomes perfectly separated, while Subgroup B goes from one misclassification to three.

The choice among these approaches requires stakeholder input about which trade-offs are acceptable in the specific application context. Engineers present these trade-offs effectively by making them explicit and quantifiable.

 Checkpoint 15.2: Fairness criteria

Fairness is not a single metric; it is a constrained design choice.

- ☐ **Metric definitions:** can you distinguish demographic parity, equal opportunity, equalized odds, and calibration in terms of which rates must match across groups?
- ☐ **Impossibility trade-off:** can you explain (at a high level) why base-rate differences make it impossible to satisfy all fairness criteria simultaneously?
- ☐ **Systems interpretation:** given the confusion matrices in table 15.6 and table 15.7, can you identify which disparity matters operationally (TPR vs. FPR vs. approval rate) and what kind of harm it represents?

- **Engineering decision:** for a concrete high-stakes domain (credit, hiring, criminal justice), can you justify which fairness constraint you would prioritize and why?

15.3.3.2 Quantifying the fairness-accuracy trade-off

The impossibility result in Kleinberg et al. (2017) establishes that fairness criteria can conflict, and figure 15.1 turns that conflict into an engineering trade-off. However, knowing the trade-off exists is insufficient: engineers must quantify the practical cost of fairness constraints to inform stakeholder decisions. A compact hiring scenario makes that cost concrete, distinct from the loan approval example in section 15.3.3.1 and with different disparity magnitudes to illustrate a different point.

Napkin Math 15.2: The price of fairness

Problem: Stakeholders demand elimination of a 20 percentage-point true positive rate (TPR) disparity in a hiring model. What is the “Price of Fairness” in terms of hiring quality?

Physics: TPRs can be equalized by adjusting the classification threshold (γ_{cls}) for the disadvantaged group.

- **Original state:** Group A (TPR = 90 percent), Group B (TPR = 70 percent). Aggregate Accuracy = 85 percent.
- **Intervention:** Lower $\gamma_{cls, g=B}$ until $TPR_{g=B} = 90$ percent.
- **The cost:** Lowering the threshold increases false positives (hiring candidates who do not meet the bar).

Math:

1. Under the scenario’s assumed threshold response, closing the 20 percentage-point TPR gap produces a 15 percentage-point increase in false positives for the disadvantaged group.
2. If the positive base rate is 20 percent, the value of a successful hire is \$100,000, and the cost of a bad hire is \$50,000, both sides of the intervention must be counted:
 - $\Delta \text{Utility} = \Delta \text{TPR} \times \text{Base Rate} \times \text{Hire Value} - \Delta \text{FPR} \times (1 - \text{Base Rate}) \times \text{Bad Hire Cost}$.
 - The added true-positive value is \$4,000 per Group B applicant, while the added false-positive cost is \$6,000, for a net loss of \$2,000.
 - This is 20 percent of Group B’s baseline utility. With Group B at 30 percent of applicants, the population-average loss is \$600 per applicant; an aggregate percentage also requires the other group’s baseline utility.

Systems insight: The “Price of Fairness” in this scenario is a 20 percent within-group utility loss, not a derivable aggregate percentage. The loss is not automatic: when a TPR gap reflects a miscalibrated threshold, closing it can raise net utility. It appears when the baseline threshold is near the utility optimum and marginal admissions skew unqualified. We owe stakeholders both the Pareto frontier and the assumptions that produced it.

The calculation gives stakeholders a way to choose a point on the fairness-accuracy frontier, but it still does not explain any particular decision. When a loan applicant receives a rejection, stating that “the model’s true positive rate for this demographic group is 60 percent compared to 90 percent for other groups” provides no actionable information. The applicant needs to know *why* the application was rejected and *what* could be changed. These questions require explainability, which is the ability to articulate which input features drove specific predictions.

15.3.4 Explainability requirements

A loan applicant denied credit by an algorithmic system may be entitled under applicable law to reasons or specific adverse-action factors, not merely aggregate statistics. **Explainability**¹⁴ supports

¹⁴ | **Explainability and interpretability:** *Interpretability* usually describes how readily a person can understand a model or its behavior, often through a constrained structure such as a short rule list or sparse linear model. *Explainability* is broader and includes post-hoc methods such as LIME and SHAP. Neither property is automatic: a linear model with opaque features can be difficult to interpret, and a feature attribution is not a causal account of a decision. The systems implication is that intrinsic constraints affect model selection, while post-hoc methods add a computation and validation path whose cost depends on the method, model, and serving workflow. GDPR access and automated-decision provisions use the phrase “meaningful information about the logic involved” for covered processing (European Parliament and Council of the European Union 2016), leaving the appropriate technical approach dependent on the decision and governing law.

this capability: it enables human oversight of automated decisions, supports debugging when problems emerge, and can satisfy applicable requirements for decision transparency.

The level of explainability required varies by application context and regulatory environment. Table 15.9 maps common deployment scenarios to their explainability needs.

Table 15.9: Explainability Considerations by Domain: Requirements vary by decision, actor, jurisdiction, and governing rules. The engineering challenge is matching explanation mechanisms to applicable requirements while managing risks such as gaming.

Application Domain	Explainability Level	Typical Requirements
Credit decisions	Specific reasons for covered actions	Principal reasons may have to be disclosed to the applicant
Medical diagnosis	Decision support	Support clinical review and applicable documentation
Content moderation	Context-dependent	Support notices or appeals where required
Recommendation	Context-dependent	Provide transparency appropriate to the use and governing rules
Fraud detection	Controlled disclosure	Balance applicable notice duties with adversarial-gaming risk

Engineering teams should select explainability approaches based on these domain requirements:

- **Post-hoc explanation methods:** Methods such as SHapley Additive exPlanations (SHAP) and Local Interpretable Model-agnostic Explanations (LIME) generate feature importance scores for individual predictions without requiring model architecture changes.¹⁵
- **Inherently interpretable models:** Expose decision structure through suitably constrained linear models, decision trees, or rule lists, but may sacrifice predictive performance. Attention weights alone are not necessarily faithful explanations of model behavior.
- **Concept-based explanations:** Map model behavior to human-understandable concepts rather than raw features.

The choice involves trade-offs between explanation fidelity, computational cost, and model flexibility.

 **Example 15.2: The hospital shortcut**

Scenario: Pneumonia detection vision models trained at Mount Sinai achieved high internal accuracy but failed on external hospital datasets (Zech et al. 2018).

Diagnosis: The neural network learned shortcut features (hospital-specific scanner artifacts and text tags) rather than true biological lung pathology.

Systems lesson: Models exploit the path of least resistance in feature spaces. External validation is essential for detecting shortcut learning; saliency maps can provide supporting evidence but do not establish that the model learned the intended mechanism.

The hospital shortcut shows why interpretability is a systems requirement rather than presentation polish: teams need enough visibility to investigate shortcuts before deployment. Figure 15.4 arranges the resulting trade-offs along a single axis. On the left side, suitably constrained decision trees and linear models can offer direct auditability, although feature design and model size still matter. On the right side, deep neural networks and convolutional architectures can provide greater capacity for complex tasks but resist direct human inspection, motivating post-hoc summaries such as LIME or SHAP that must themselves be validated.

The choice depends on the application’s accountability requirements. Adverse-action laws require accurate, specific reasons for covered credit decisions, but they do not mandate a particular model architecture. Other applications may face different transparency, safety, or contestability duties. The spectrum does not imply “simple is always better,” because a highly interpretable model that makes inaccurate predictions may also cause harm. The engineering challenge is selecting a model and explanation process that meet the application’s predictive and accountability requirements.

Some explainability and transparency requirements carry the force of law. The EU AI Act, which entered into force on August 1, 2024 and applies in phases, imposes documentation, transparency,

15 | **LIME (local interpretable model-agnostic explanations) and SHAP (SHapley additive explanations):** LIME (Ribeiro et al. 2016) fits a local interpretable surrogate around each prediction; SHAP (Lundberg and Lee 2017) adapts Shapley values from cooperative game theory to compute feature contributions under a unified additive framework. Exact Shapley-value computation can be expensive, so practical SHAP implementations rely on approximations or model-specific algorithms. The systems trade-off is that explanation fidelity, latency, and implementation complexity must be budgeted explicitly rather than treated as free.




Figure 15.4: Model Interpretability Spectrum: A horizontal continuum groups decision trees, linear regression, and logistic regression as intrinsically interpretable, while random forests, neural networks, and CNNs/transformers require post-hoc explanation. The grouping frames model selection as an auditability trade-off rather than a universal ranking of model quality.

human-oversight, and risk-management obligations for covered high-risk systems (European Parliament and Council of the European Union 2024). Regulation B requires specific principal reasons for covered adverse actions, and its official interpretation states that disclosed reasons must accurately describe factors actually considered or scored (12 C.F.R. § 1002.9; official interpretation). The applicable technical mechanism depends on the system, decision, jurisdiction, and legal obligation.

15.3.5 The regulatory landscape

Regulation changes responsible engineering from a best-practice argument into architecture constraints. Imagine a credit model denies an applicant and the applicant asks why, contests the decision, and later requests access to the data used about her. The system must do more than report an accuracy score. It must produce an explanation tied to the specific decision, preserve the model and data lineage that led to that output, route the dispute to a substantive human review path, retain audit logs, and support deletion or access workflows where data rights apply. Responsible engineering now operates within explicit regulatory frameworks that turn transparency, oversight, and accountability into technical requirements.

Regulation first enters the architecture through risk classification. The EU AI Act establishes a comprehensive framework, classifying AI systems by risk level and mandating requirements accordingly.¹⁶ The Act entered into force on August 1, 2024 and applies in phases: prohibited-practice rules began applying in 2025, while high-risk and other operator obligations phase in by system category and implementation guidance. Article 99 sets maximum fines of EUR 35 million or 7 percent of global turnover for prohibited AI practices, while many other operator obligations are capped at EUR 15 million or 3 percent (European Parliament and Council of the European Union 2024).

For engineers, the important point is not the fine schedule but the capabilities the law demands. Covered high-risk systems¹⁷ must implement risk management, data governance, technical documentation, transparency, human oversight, and accuracy, robustness, and security requirements. A covered credit-decision system therefore needs auditability from inception: model versions, training data provenance, validation evidence, human-oversight design, logging, and postdeployment monitoring must be part of the architecture rather than documents assembled after launch.

Contestability adds a second architectural requirement. GDPR moves the same applicant workflow into data-subject rights. Article 22 grants EU data subjects the right not to be subject to decisions based solely on automated processing that produce legal or similarly significant effects, subject to specified exceptions and safeguards.¹⁸ Article 15(1)(h) separately gives data subjects access to meaningful information about the logic involved in automated decision-making referred to in Article 22. Engineering teams should determine which decisions are covered and design the required explanation and human-review capabilities accordingly. Where a substantive review path is required, it must be operationally staffed and supported by summaries, provenance, and audit tools.

US sectoral law reaches similar capabilities through domain-specific evidence requirements. These regulations are less unified than the EU AI Act, but they can impose related engineering duties. In the credit example, the Equal Credit Opportunity Act (ECOA) and its implementing Regulation B require specific principal reasons for covered adverse actions, and those reasons must accurately

¹⁶ | **EU AI Act (Regulation 2024/1689):** The first comprehensive AI legal framework, defining four risk tiers with penalties that vary by infringement category; prohibited AI-practice violations can reach EUR 35 million or 7 percent of global turnover; many other obligations, including many high-risk operator obligations, are capped at EUR 15 million or 3 percent. The Act has extraterritorial reach: non-EU organizations may need to comply when they place systems on the EU market or when system outputs are used in the EU. Systems engineering implications are concrete: high-risk AI requires logging infrastructure for audit trails, human oversight mechanisms built into the architecture, and CE marking—all capabilities that must be designed in from inception, not retrofitted after deployment.

¹⁷ | **High-risk AI (EU AI Act Annex III):** Annex III enumerates specified use cases in areas including biometrics, critical infrastructure, education, employment, essential services, law enforcement, migration, and justice. Classification depends on intended use and the Article 6 criteria, including exclusions for some systems that do not materially influence decisions or pose significant risk; profiling systems within Annex III remain high-risk. Model architecture alone does not determine classification.

18 | GDPR (General Data Protection Regulation) articles 15 and 22: Article 22 restricts certain solely automated decisions with legal or similarly significant effects and, for specified exceptions, requires safeguards including human intervention, the ability to express a point of view, and the ability to contest the decision; article 15(1)(h) contains the access right to meaningful information about the logic involved in such automated decision-making (European Parliament and Council of the European Union 2016). The European Data Protection Board's guidance emphasizes that required human oversight must be substantive and not merely a "rubber-stamping" exercise (European Data Protection Board 2018). If 0.1 percent of 1M daily decisions are appealed or escalated, the system must handle 1,000 cases/day; this is a workload assumption, not a model-error rate.

19 | HIPAA (Health Insurance Portability and Accountability Act): Enacted in 1996, with Privacy Rule and Security Rule requirements establishing standards for protected health information (United States Congress 1996; U.S. Department of Health and Human Services 2003, 2005); PHI may be used for training under applicable authorization or another permitted pathway, such as an IRB or Privacy Board waiver; de-identified data and limited data sets with data-use agreements provide additional pathways. Model outputs may remain PHI when they identify or can reasonably identify individuals, and security-rule documentation retention must be reflected in audit and evidence design. Civil money penalties are tiered and inflation-adjusted by regulation (U.S. Department of Health and Human Services 2026).

describe factors actually considered or scored (12 C.F.R. § 1002.9; official interpretation). If consumer-report information or credit scores influence the decision, the Fair Credit Reporting Act (FCRA) adds notice obligations; if the same scoring machinery is used for housing, the Fair Housing Act (FHA) adds a discrimination-prohibition constraint. The technical consequence is practical rather than abstract: covered systems need evidence and controls that support the applicable reasons, notices, review, and nondiscrimination duties.



Checkpoint 15.3: Ethical deployment

Deployment is where safeguards must become operational.

Safety net

- ☐ **Rollback:** verify that the previous model can be restored within the incident-response target.
- ☐ **Human-in-the-loop:** is there a path for human review of low-confidence predictions?

Monitoring plan

- ☐ **Silent failure:** specify the subgroup metrics that will expose post-deployment bias.

Healthcare regulations, including the Health Insurance Portability and Accountability Act (HIPAA)¹⁹ and Food and Drug Administration (FDA) guidance, impose the same pattern with different artifacts: protected-health-information controls, validation records, audit logs, and incident response. Employment systems likewise require evidence that automated screening does not reproduce discriminatory hiring practices. Across domains, the task is to translate each obligation into a concrete capability: explanation, human review, lineage, access control, deletion, monitoring, or incident response. The deployment checkpoint is therefore not a US-sectoral checklist; it is the common production contract implied by the regulatory landscape.

The engineering response to these regulatory requirements is proactive architectural design. Teams that build documentation, monitoring, explainability, and human oversight into systems from inception demonstrate compliance efficiently. Teams that must retrofit these capabilities face expensive redesign or deployment constraints. The foundation established here, that responsibility is an engineering requirement rather than a legal afterthought, enables more targeted compliance strategies as regulatory frameworks mature. Yet regulatory readiness still covers only the planned path; even well-designed systems can fail, making incident response preparation essential.

15.3.6 Monitoring and incident response

Zillow reported a \$304 million²⁰ Q3 2021 Homes-segment inventory write-down after buying homes at prices above revised estimates of future selling prices (Zillow Group 2021). A systems diagnosis can interpret the failure as a combination of forecasting uncertainty, distribution shift, operational capacity limits, and insufficient circuit breakers. Planning for system failures before they occur is a core responsible engineering practice. Incident response and monitoring require preparation before the system fails. Building on the incident severity classification and response framework from section 14.5.4.1, table 15.10 adapts that general framework to responsible deployment, where detection must surface fairness violations and demographic-slice degradation alongside ordinary outages. The five components are largely the standard incident-response arc; what makes the table actionable is its last column. The requirements state what each component must do, but the predeployment-verification column states what must be proven before launch: an alert threshold that has been tested, a rollback path that has been exercised, a contact list that is current. A requirement without a verified control is an intention, not a safeguard, so this column is the gate that decides whether the system is ready to deploy.

Table 15.10: Incident Response Framework: Systematic preparation for ML system failures requires five distinct components. Detection identifies anomalies through specialized monitoring; assessment evaluates scope using severity classifications; mitigation reduces harm through tested rollback procedures; communication notifies stakeholders through preapproved channels; remediation implements permanent fixes through root cause analysis. Each component requires both operational requirements and predeployment verification.

Component	Requirements	Predeployment Verification
Detection	Monitoring systems that identify anomalies, degraded performance, and fairness violations	Alert thresholds tested, on-call rotation established, escalation paths documented
Assessment	Procedures for evaluating incident scope and severity	Severity classification defined, impact assessment templates prepared
Mitigation	Technical capabilities to reduce harm while investigation proceeds	Rollback procedures tested, fallback systems operational, kill switches functional
Communication	Protocols for stakeholder notification	Contact lists current, message templates prepared, approval chains defined
Remediation	Processes for permanent fixes and system improvements	Root cause analysis procedures, change management integration

ML systems create unique maintenance challenges and technical debt (Sculley et al. 2015). Models degrade silently, dependencies shift unexpectedly, and feedback loops amplify small problems into large ones. Incident response planning must account for these ML-specific failure modes, and effective response depends on continuous monitoring infrastructure that detects problems in the first place. The monitoring infrastructure from chapter 14 provides the foundation for responsible system operation, extending traditional operational metrics to include outcome quality measures.

Responsible monitoring extends along several interconnected dimensions:

- **Performance stability tracking:** Detects gradual prediction quality degradation that might not trigger immediate alerts. Slow accuracy decay that accumulates over weeks can be particularly dangerous because it evades threshold-based alarms.
- **Subgroup parity monitoring:** Adds a fairness lens to temporal tracking, comparing error rates across demographic groups to detect emerging disparities before they cause significant harm.
- **Input distribution monitoring:** Can flag population shifts and potential adversarial manipulation before observed outcome metrics move.
- **Outcome monitoring:** Validates whether predictions translate to intended real-world results, not merely whether model scores remain stable.
- **User feedback systems:** Surface complaints and corrections that reveal problems invisible to any automated metric, including harms that only affected users can articulate.

Together, these dimensions connect model-level metrics, data-layer shifts, real-world outcomes, and human reports into one monitoring surface.

Effective monitoring requires both data collection infrastructure and disciplined review processes. Dashboards that no one examines provide no protection, so engineering teams must establish regular review cadences with clear ownership and escalation procedures.

The frameworks established in this section address one dimension of responsible engineering: ensuring systems work fairly and reliably across user populations. Fairness is not the only cost that conventional engineering metrics overlook. Every model training run, every inference request, every monitoring dashboard consumes electricity that translates into carbon emissions and dollar costs. A system can be perfectly fair across demographic groups while consuming orders of magnitude more resources than the task requires, harming not specific user populations but the broader environment and the organizations paying the bills. Responsible engineering must therefore extend beyond *who* the system serves to encompass *what it costs* to serve them.

20 | **Zillow's D·A·M (data · algorithm · machine) failure:** Zillow's 2021 write-down is a useful systems case because the documented business failure combined forecast uncertainty with operational execution (Zillow Group 2021). A D·A·M diagnosis interprets the data axis as the mismatch between historical home-sale data and pandemic-era price volatility, the algorithm axis as the difficulty of pricing homes with reliable uncertainty estimates, and the machine axis as an automated iBuying pipeline that needed stronger capacity limits and circuit breakers. This is an engineering interpretation of Zillow's public disclosure, not a claim that Zillow identified one root technical cause.

15.4 Environmental and Cost Awareness

In 2019, researchers estimated that development-scale training and architecture search for a large Natural Language Processing (NLP) model could emit as much carbon as five cars over their entire lifetimes (Strubell et al. 2019). Later analysis showed that the most quoted architecture-search estimate was highly sensitive to proxy-task, hardware, data-center efficiency, and grid carbon-intensity assumptions (Patterson et al. 2021). The correction sharpened the responsible-engineering point rather than weakening it: training runs consume megawatt-hours of electricity, inference at scale multiplies per-request inefficiencies into measurable environmental impact, and resource-intensive models exclude organizations that lack large compute budgets. The optimization techniques developed in chapter 10, chapter 11, and chapter 12 therefore serve double duty as instruments of responsible engineering, connecting computational efficiency to environmental sustainability, economic accessibility, and long-term scalability.

15.4.1 Efficiency as responsibility

Training a single large language model consumes thousands of GPU hours and energy measured in megawatt-hours. Much of this expense, however, is not intrinsic to the learning task but represents *accidental complexity*: training from scratch when fine-tuning would suffice, using larger models than tasks require, and running hyperparameter searches that explore redundant configurations. Computational cost depends on engineering choices as well as model physics. Green AI treats that efficiency as a primary metric rather than an afterthought.²¹

Resource efficiency and responsible engineering are directly linked through three interconnected channels:

- **Environmental impact:** More computation can increase energy use and emissions, but the relationship also depends on hardware utilization, power, runtime, facility overhead, and grid intensity. Efficiency techniques reduce environmental impact when they lower measured lifecycle energy for the required workload and quality target.
- **Accessibility:** Resource-efficient models can broaden deployment options by running on less expensive hardware. A quantized model that runs on a smartphone can reduce dependence on connectivity and recurring cloud API costs.
- **Sustainability at scale:** Systems serving millions of users multiply per-request resource demand. A latency reduction saves accelerator time only when it reduces service demand rather than shifting work or idle time; annual savings depend on traffic, batching, concurrency, and utilization.

These channels make optimization a responsibility requirement rather than a narrow performance exercise.

The optimization techniques can directly serve responsibility goals. Quantization, pruning, knowledge distillation, and hardware acceleration can reduce model size, executed work, latency, or energy, but the realized benefit and quality impact depend on the model, workload, runtime, and hardware.

Responsible engineers apply these techniques as design requirements, not afterthoughts. The question shifts from maximizing accuracy alone to maximizing accuracy within efficiency constraints.

15.4.2 Efficiency engineering in practice

Acknowledging that efficiency matters is the easy part; the harder engineering challenge is translating that principle into measurable targets. The goal is selecting the smallest model that meets task requirements, then applying methodical optimization to reduce resource consumption further. Edge deployment scenarios make these constraints concrete because they impose hard physical limits that cannot be negotiated away.

Edge deployment scenarios make efficiency requirements concrete. In this illustrative scenario, a wearable device has a 500 mW power budget and must run inference continuously for 24 hours on a small battery. Table 15.11 compares assumed constraints across four deployment contexts, from smartphones with 5 W budgets to IoT sensors operating at 100 mW.

²¹ | **Green AI:** Schwartz et al. (2020) contrasted “Red AI” (performance at any cost) with “Green AI” (efficiency as primary metric). The compute-growth anchor comes from AI and Compute’s 2012–2018 trend analysis, which reported a 300,000× increase in compute used in the largest AI training runs (Amodei and Hernandez 2018b). The Green AI proposal—reporting FLOPs alongside accuracy for every published result—reframes efficiency from an engineering preference into a scientific reporting obligation, making the resource cost of marginal accuracy gains visible and comparable across research groups.

Table 15.11: Illustrative Edge Deployment Constraints: The power and latency values are scenario assumptions for comparing deployment contexts; actual budgets depend on hardware, workload, and product requirements.

Deployment Context	Power Budget	Latency Requirement	Typical Use Cases
Smartphone	5 W	100 ms	Photo enhancement, voice assistants
IoT Sensor	100 mW	1 second	Anomaly detection, environmental monitoring
Embedded Camera	1 W	30 FPS (33 ms)	Real-time object detection, surveillance
Wearable Device	500 mW	500 ms	Health monitoring, activity recognition

The scenarios in table 15.12 provide guidance for efficiency optimization. Techniques that enable deployment on power-constrained platforms can reduce environmental impact per inference when they lower measured energy without shifting costs elsewhere. Financial savings also depend on provisioning, utilization, pricing, and traffic.

Table 15.12: Model Efficiency Comparison: Illustrative model profiles under the stated deployment assumptions. Model selection must account for accuracy, latency, power, memory, runtime support, and lifecycle impact; parameter count alone does not determine cost or environmental impact.

Model	Parameters	Inference Power	Latency	Fits Smartphone?	Fits IoT?
MobileNetV2	3.5M	1.2 W	40 ms	Yes	No
EfficientNet-B0	5.3M	1.8 W	65 ms	Yes	No
ResNet-50	25.6M	4.5 W	180 ms	No	No
TinyML Model	200K	50 mW	200 ms	No	Yes

For the wearable budget in table 15.11, the TinyML model leaves a 10× power margin, while MobileNetV2 exceeds the same power budget by 2.4× before accounting for sustained thermals. Fitting the device envelope is necessary but not sufficient: per-inference power compounds into lifetime serving cost once the model runs continuously at production scale.

15.4.3 Total cost of ownership

A team spends \$3,200 training a recommendation model and celebrates the modest cost. Six months later, they discover they are spending \$500,000 per year serving it. The surprise illustrates how total cost of ownership²² can be dominated by recurring inference at high traffic. Other systems may be dominated by training, data, staffing, or idle capacity, so measurement determines where optimization should focus.

Consider a concrete example of a recommendation system serving 10M users daily. Training costs appear considerable: data preparation consumes 100 GPU-hours at approximately \$4/hour (\$400), hyperparameter search across multiple configurations requires 500 GPU-hours (\$2,000), and the final training run uses 200 GPU-hours (\$800). Total training cost reaches approximately \$3,200.

Inference costs dominate in this scenario. With 10M users each receiving 20 recommendations per day, the system serves 200M inferences daily. Treating 10 milliseconds per inference as unbatched, non-overlapped dedicated GPU service demand yields approximately 23.1 GPUs running continuously. At \$2.50/GPU-hour, the scenario's annual GPU cost reaches \$506,944.

Over a three-year operational period, quarterly retraining produces total training costs of approximately \$38,400, while inference costs over the same period total \$1.5M. The resulting 40:1 ratio is specific to these assumptions, but it tells us where to direct optimization effort in this scenario: inference latency and serving efficiency.

Per-query optimization becomes important when serving billions of requests. Shaving a fifth off the per-query accelerator service demand can reduce required hardware when other workload assumptions remain fixed. Hardware selection among CPUs, GPUs, and Tensor Processing Units changes costs and carbon footprint in workload- and deployment-dependent ways. Model compression through quantization and pruning can reduce high-volume inference costs when it lowers measured service demand without unacceptable quality loss.

22 | TCO (total cost of ownership): ML TCO includes labeling, monitoring, retraining, remediation, energy, audits, and compliance. At sufficient traffic, recurring inference can dominate a one-time training bill; utilization and update cadence determine the balance.



At production scale, recurring inference can dominate lifetime cost.

Total cost of ownership (TCO) encompasses additional dimensions beyond computation. Operational costs include monitoring, maintenance, retraining, and incident response, all of which scale with system complexity and the rate of distribution shift in the application domain. Opportunity costs reflect that resources consumed by ML systems cannot be used for other purposes. Wasteful resource consumption in one project constrains what other projects can attempt.

Engineers should evaluate return on investment (ROI): whether the value an ML system delivers justifies its resource consumption. A recommendation system that increases engagement by 1 percent might not justify millions of dollars in computational costs, while a medical diagnosis system that saves lives does. Explicit trade-offs enable responsible resource allocation.²³

15.4.3.1 TCO calculation methodology

Quantifying operational carbon impact requires measured or estimated energy use, facility overhead, and applicable grid carbon intensity. Engineers can estimate three-year total cost of ownership using a structured approach that separates training, inference, and operational costs into ledgers. Training is usually a one-time or periodic expense, inference recurs with every user request, and operations accumulate through monitoring, retraining, and incident response. The ledger methodology applies that structure to the recommendation system example in this section.

 Napkin Math 15.3: The carbon cost of compute

Problem: The TCO ledgers must convert “compute hours” into “kg CO₂eq”. What carbon factor does one GPU-hour carry under the scenario baseline?

Variables:

- **Power:** 400 W per GPU (scenario baseline).
- **Intensity:** 0.4 kg/kWh CO₂eq (rounded grid baseline).

Math: Equation 15.2 captures the standard conversion:

$$\text{Carbon} = \text{Energy (kWh)} \times \text{Carbon Intensity (kg/kWh)} \tag{15.2}$$

Applying the baseline assumptions:

$$(0.4 \text{ kW} \times 1 \text{ hour}) \times 0.4 \text{ kg/kWh} = 0.16 \text{ kg CO}_2\text{eq per GPU-hour}$$

Systems insight: This conversion factor lets the ledgers track “Carbon Cost” alongside “Dollar Cost”, making emissions a first-class engineering metric across the downstream TCO tables.

Training costs Training costs include both initial development and ongoing retraining. Table 15.13 breaks down these costs, showing how quarterly retraining cycles accumulate over a three-year operational period.

Table 15.13: Training Cost Calculation: Training costs accumulate through initial development (\$3,200 per cycle) and quarterly retraining over a three-year operational period. Data preparation, hyperparameter search, and final training consume GPU hours at \$4/hour, totaling \$38,400 across 12 training cycles. Despite appearing substantial, training represents only 1.9 percent of total cost of ownership.

Cost Component	Calculation	Financial Cost	Carbon (kg CO ₂)
Initial data preparation	hours × rate	100 GPU-hr × \$4 = \$400	16 kg
Hyperparameter search	experiments × cost/experiment	50 × \$40 = \$2,000	80 kg
Final training	hours × rate	200 GPU-hr × \$4 = \$800	32 kg
Subtotal per training cycle		\$3,200	128 kg
Retraining frequency	cycles/year × years	4/year × 3 years = 12	same multiplier (12 cycles)
Total training cost	subtotal × cycles	\$38,400	1,536 kg

Inference costs Table 15.14 walks the conversion chain that turns traffic into cost: daily queries become GPU-seconds, GPU-seconds become GPU-hours, and GPU-hours convert into both dollars

23 | **ML ROI (return on investment):** Deployment-to-training cost ratios vary with traffic, utilization, hardware, update cadence, staffing, and incident burden. The appropriate model choice depends on measured lifecycle cost and task value rather than a fixed ratio.

and carbon. Carbon attaches only once the workload is expressed in GPU-hours, so the query and GPU-second rows leave the carbon column blank by design rather than omitting data. Following the chain to the bottom row supplies the inference term for the total-cost comparison in table 15.16; dominance follows only when that term is compared with training and operations.

Table 15.14: Illustrative Inference Cost Calculation: In this unbatched, non-overlapped service model, 200M daily queries at 10 ms of dedicated accelerator demand require 556 GPU-hr daily, totaling \$507K annually and \$1.52M over three years. Real utilization, batching, and concurrency change this conversion.

Cost Component	Calculation	Financial Cost	Carbon (kg CO ₂)
Daily queries	users × queries/user	10M × 20 = 200M	-
GPU-seconds/day	queries × service demand	200M × 0.01 s = 2M sec	-
GPU-hours/day	seconds ÷ SEC_PER_HOUR	556 GPU-hr	88.9 kg
Annual GPU cost	hours × 365 × rate	556 × 365 × \$2.50 = \$507K	32,444.4 kg
3-year inference cost	annual × 3	\$1.52M	97,333.3 kg

Operational costs Operational costs encompass infrastructure, personnel, and incident response. ML systems generate operational burdens that traditional software does not: “incident response” frequently means debugging silent failures (data drift, feature corruption, or distribution shifts) rather than binary service outages, and “monitoring infrastructure” must continuously track statistical anomalies in model predictions across demographic slices, not merely service availability. Table 15.15 itemizes these ongoing expenses, which often surprise teams focused primarily on compute costs.

Table 15.15: Operational Cost Calculation: Illustrative operational costs include monitoring infrastructure (\$50K/year), on-call engineering at 0.5 FTE (\$100K/year), and incident response reserves (\$20K/year). The \$510K three-year total represents 24.6 percent of this scenario’s TCO. Actual staffing and incident costs depend on system scale, risk, organizational structure, and support model.

Cost Component	Annual Estimate	3-Year Total
Monitoring infrastructure	\$50K	\$150K
On-call engineering (0.5 FTE)	\$100K	\$300K
Incident response (estimated)	\$20K	\$60K
Total operational		\$510K

The stark breakdown in table 15.16 answers where the money goes: inference at 73.5 percent, operations at 24.6 percent, and training at only 1.9 percent.

Table 15.16: Total Cost of Ownership Summary: In this illustrative service-demand scenario, the inference-to-training cost ratio is 40:1. A 20 percent reduction in per-query accelerator demand saves about \$304K and 19 t of modeled accelerator CO₂ under the same assumptions. The carbon subtotal excludes operations, facility overhead, and embodied impacts.

Category	3-Year Cost	Percentage	Modeled Accelerator Carbon
Training	\$38K	1.9%	1.5 t
Inference	\$1.52M	73.5%	97.3 t
Operations	\$510K	24.6%	Not estimated
Total TCO	\$2.07M	100%	~98.9 t modeled subtotal

Those proportions turn efficiency from a tuning preference into a responsibility check.



Checkpoint 15.4: Efficiency as responsibility

Total cost of ownership reveals where responsible optimization has the most leverage.

- ☐ **Inference dominance:** under this scenario's assumptions, can you explain why a 20 percent accelerator-service-demand reduction delivers more savings than a 50 percent training time reduction?
- ☐ **Carbon accounting:** can you convert GPU-hours into kg CO₂eq using power draw and carbon intensity, and explain how workload design, hardware, scheduling, and region affect the result?
- ☐ **Sufficiency test:** for a given ML system, can you justify that the model size is appropriate for the task, or identify where a simpler model would deliver comparable accuracy at a fraction of the cost?

15.4.3.2 Environmental impact

The TCO analysis in the preceding section captures costs that appear on invoices, but computational resources carry costs that no invoice reflects. Operational emissions depend on measured energy, facility power usage effectiveness (PUE), and regional grid carbon intensity (gCO₂e/kWh); embodied manufacturing impacts require separate accounting. A first-order operational estimate is $\text{Energy}_{\text{hardware}} \times \text{PUE} \times \text{Carbon Intensity}$. Per-request energy is measured in joules per inference, whereas joules per FLOP characterizes hardware work efficiency. An optimization that reduces TCO may not reduce energy or emissions if it changes utilization, hardware, or workload placement. Data-center electricity use makes workload design, hardware, cloud region, and timing part of responsible engineering (Henderson et al. 2020). The magnitude becomes clearer in a scale calculation for training a large foundation model.

Efficiency optimization and environmental responsibility align when measured lifecycle energy and emissions fall for the required workload and quality target. More granular carbon accounting methodologies build on this foundation: lifecycle assessment tracks impacts across the system's full life, scope 1/2/3 emissions separate direct emissions, purchased electricity, and supply-chain or use-phase emissions, and carbon-aware scheduling shifts work toward lower-carbon times or regions.

The same physical quantities that govern performance also affect responsibility. Data movement and computation contribute to chip-level energy, while data-center emissions additionally depend on utilization, facility overhead, grid intensity, and embodied impacts. Pareto analysis applies to both accuracy-fairness and accuracy-latency objectives, but reweighting an objective can improve multiple metrics when the current solution is dominated. *Responsible engineering extends the constrained optimization problem this book has been teaching to objectives that include societal impact alongside throughput and latency.*



Napkin Math 15.4: The carbon cost of scale

Problem: A foundation model is being trained at the scale of GPT-3, consuming 1,287 MWh of electricity (Patterson et al. 2021). What is the operational carbon impact under the stated grid-intensity assumption?

Math:

1. **Energy consumption:** 1,287 MWh = 1,287,000 kWh.
2. **Carbon intensity:** This scenario uses ≈ 429 g/kWh CO₂ (0.429 kg/kWh) (Patterson et al. 2021).
3. **Total emissions:** 1,287,000 kWh $\times$ 0.429 kg/kWh = 552,123 kg CO₂ (552 t).
4. **Comparison:** The scenario assumes 4.6 t CO₂ per passenger-car year.

Systems insight: Under this training-energy scenario, the operational emissions equal the assumed annual emissions of 120 passenger cars. A 1 percent reduction in total training energy

Train

100x+

Car-year

This training scenario exceeds 100 passenger-car years of emissions.

under the same grid mix would remove the equivalent of about 1.2 cars taken off the road for a year from the calculation.

The checklists, fairness metrics, explainability mechanisms, and efficiency analyses developed in previous sections tell engineering teams *what to measure* and *how to act*. A natural follow-up concern is *what* infrastructure records those answers, supports audits, and routes violations to the appropriate intervention rather than relying on human memory. The answer lies in **data governance**—the engineering discipline that transforms policy intentions into enforceable technical controls.

15.5 Data Governance and Compliance

Recommendation and targeting models can consume user data at enormous scale, which means a governance failure in the data pipeline is simultaneously a failure in the ML system’s data ingestion and validation infrastructure. Governance is the enforcement layer that makes accountability possible: fairness metrics, model cards, and impact assessments only matter at scale if the system can establish what data it used, who accessed it, which legal basis allowed processing, and whether applicable deletion or contestability rights were honored.

The Meta fines make the governance question concrete: a system must establish a lawful basis for processing as well as protect the data. In January 2023, the Irish Data Protection Commission issued separate EUR 210M and EUR 180M fines (totaling EUR 390M) against Meta Ireland. The decisions concerned unlawful reliance on contractual necessity for personalized advertising together with transparency and fairness violations, not a data breach (Data Protection Commission 2023).

The storage architectures examined in chapter 4 are governance enforcement mechanisms that determine who accesses data, how usage is tracked, and whether systems comply with regulatory requirements. Every architectural decision, from acquisition strategies through processing pipelines to storage design, carries governance implications that manifest when systems face regulatory audits, privacy violations, or ethical challenges. Data governance turns abstract policy into concrete engineering through access controls on training data, audit infrastructure that records data access, privacy-preserving training techniques, and lineage systems that trace raw audio recordings to production models. Because those obligations bind architecture rather than paperwork, they are worth stating as a principle in their own right, one local to this chapter rather than one of the part-opening principles that frame the volume.

Principle: Compliance as engineering constraint

Invariant: Data governance obligations are engineering requirements, not optional policy overlays.

Implication: Systems that process regulated data must map applicable duties to technical and organizational controls. Access control, erasure, contestability, audit, and lineage may be required or useful depending on the jurisdiction, data, actor, and use case. The General Data Protection Regulation (GDPR), California Consumer Privacy Act (CCPA), Brazil’s General Data Protection Law (LGPD, *Lei Geral de Proteção de Dados*), and China’s Personal Information Protection Law (PIPL) differ in scope and duties, so no single control checklist establishes compliance across them.

Data governance encompasses four interconnected domains that make security, privacy, compliance, and lineage mutually dependent enforcement constraints:

- **Security infrastructure:** Protects data assets through identity controls, access authorization, encryption, and secure key management across the data lifecycle.
- **Privacy mechanisms:** Determine what information is exposed even to authorized users, respecting individual rights while enabling model training.
- **Compliance frameworks:** Translate jurisdiction-specific regulatory requirements into architectural constraints that shape how data flows through the system.

- **Lineage and audit systems:** Create the accountability trails that make the first three domains verifiable. Without them, security policies, privacy guarantees, and compliance claims are difficult to demonstrate or investigate.

The lighthouse KWS system, the keyword-spotting voice assistant introduced in chapter 2 as the Tiny Constraint lighthouse, illustrates how the fairness risks identified in table 15.5 intensify at the governance level. A deployed architecture may process audio locally in users' homes, maintain authorized voice-pattern histories, and cache models derived from population-wide training data. These capabilities create governance obligations around consent management, data minimization, access auditing, and deletion rights.

For the KWS system, the four domains become concrete enforcement questions. Security determines who may reach data through encryption at rest and in transit, role-based access controls, and logging of policy-relevant feature-store access. Privacy determines what the system may retain through data minimization and, where appropriate, differential privacy. Compliance determines which regulatory duties constrain data flow, translating applicable requirements into erasure pipelines and consent-management APIs. Lineage and audit determine how the system reconstructs what happened through documentation of data provenance, model lineage, and decision audit trails. Figure 15.5 shows the broader operating model that supports these domains: organization, policies, data catalogs, data sourcing, data quality, data operations, data security, and shared definitions. The domains and this broader operating model must work together because a failure in any one undermines the others: encrypted data with no access controls is still vulnerable, and compliant storage without verifiable audit trails cannot survive a regulatory audit. In the context of the D·A·M taxonomy, governance provides structural support for the data axis, helping keep the system's information safe, compliant, and reliable across the data lifecycle.




Figure 15.5: Data Governance Framework: Eight operating elements surround Data Governance as a central hub: policies, organization, security, operations, data quality and master data, sourcing, catalogs, and shared analytic definitions. The circular arrangement presents governance as a coordinated operating model rather than a single control, spanning organizational functions such as catalogs, sourcing, and analytic definitions alongside the four technical domains (security, privacy, compliance, and lineage). Adapted from standard enterprise data governance frameworks.

15.5.1 Security and access control architecture

Consider a data scientist querying a feature store for training data. She can read aggregated voice features but cannot access the raw audio recordings from which they were derived. The serving pipeline can read online features for inference but cannot write to the training dataset. Neither can modify source data. The separation is intentional: it reflects a layered security architecture

where governance requirements translate into enforceable technical controls at each pipeline stage. Feature stores can implement role-based access control (RBAC) that maps organizational policies into database permissions, preventing unauthorized access. These controls operate across storage tiers: object storage like S3 enforces bucket policies, data warehouses implement column-level security that hides sensitive fields, and feature stores maintain separate read/write paths with different permission requirements.

Access control mechanisms remain incomplete without encryption, which protects data at rest and in transit but does not prevent misuse after decryption or authorized access. Training data stored in data lakes can use server-side encryption with keys managed through dedicated key management services (AWS KMS, Google Cloud KMS). Feature stores can use encryption at rest and transport encryption in transit. Lighthouse KWS edge devices can combine transport encryption for confidentiality with code signing to verify model-update integrity; both controls still require secure key management and update authorization.

Serialized weights, Open Neural Network Exchange exports, and checkpoints require the same protection as training data. Weights represent substantial intellectual property and are a vulnerability surface: an attacker who injects a backdoored model into a registry can compromise deployments that accept the artifact, an attacker who retrieves weights may probe whether particular records were in training (Shokri et al. 2017), and model-extraction attacks can steal model behavior through prediction APIs (Tramèr et al. 2016). Model registries therefore require version-pinned access controls, cryptographic signatures that verify artifact integrity before serving, and write-protected promotion gates so that only pipeline-validated checkpoints can overwrite production slots. Training pipelines themselves must be protected from data poisoning attacks, where adversarially crafted samples inserted into the training corpus cause the learned model to exhibit targeted misbehavior at inference time.

Access control and encryption establish *who* can reach data and *how* it is protected in transit and at rest. Controlling access is only half the problem: even authorized users can compromise individual privacy if the data itself is insufficiently protected.

15.5.2 Technical privacy protection methods

A data scientist with legitimate access to training data does not need, and should not see, individual user records when aggregate statistics suffice. Privacy-preserving techniques²⁴ address this gap by determining what information systems expose even to authorized users, adding a second layer of protection beyond access control. Differential privacy provides a formal bound on how much one record can affect the output distribution. Implementing differential privacy in production requires careful engineering: applying a mechanism with calibrated noise, tracking privacy budgets with a valid accountant across composed data uses, and testing the implementation. Membership inference attacks can supplement these checks but cannot establish the formal guarantee.²⁵

KWS systems face particularly acute privacy challenges because the always-listening architecture requires processing audio continuously while minimizing data retention and exposure. Production systems can address these risks through three architectural choices:

- **On-device processing:** Can keep wake word detection local, with audio transmitted only after detection under the stated design.
- **Federated learning:** Allows devices to train on local audio and share model updates rather than raw recordings; aggregation and privacy controls still matter.²⁶
- **Automatic deletion policies:** Apply stated retention periods to primary storage and serving stores. Complete deletion also requires verification across replicas, backups, caches, and derived artifacts according to the governing policy.

Together, these choices minimize what leaves the device, what the server can reconstruct, and how long sensitive traces persist.

15.5.3 Architecting for regulatory compliance

When a European user invokes an applicable right to erasure under GDPR, the voice assistant must determine which personal data and downstream artifacts are in scope and execute the appropriate workflow within the regulation's response deadlines (European Parliament and Council

24 | Privacy-preserving techniques: K-anonymity (Sweeney 2002) ensures each record is indistinguishable from at least $k - 1$ others with respect to chosen quasi-identifiers, l-diversity adds attribute variety within equivalence classes (Machanavajhala et al. 2007), and t-closeness bounds distribution distance (Li et al. 2007). These syntactic guarantees address different threats from differential privacy and do not by themselves protect against all ML-specific leakage: a model trained on de-identified data can still memorize examples or reveal membership signal under some conditions (Shokri et al. 2017). Differential privacy instead bounds the influence of one record under a specified mechanism and privacy budget, offering a semantic guarantee designed to remain meaningful in the presence of side information; the methods are therefore not a simple strength hierarchy, and the appropriate protection depends on the release, threat model, and utility requirements.


25 | Membership inference attack: The attack uses model behavior to infer whether a record appeared in training; overfitting can increase the signal (Shokri et al. 2017). In practice, higher confidence on training samples can help distinguish members from non-members. Its measured success depends on model access, data distribution, and defense assumptions. A successful empirical attack reveals practical leakage, but it does not alone prove violation of a stated (ϵ, δ) bound. Formal differential privacy comes from the mechanism and its accounting.

26 | Federated learning:

McMahan et al. (2017) introduced Federated Averaging (FedAvg), in which each device trains locally and shares model updates rather than raw records. This architecture can limit routine transfer of raw data, but it is not a privacy guarantee: model updates can leak training information through reconstruction attacks (Zhu et al. 2019). Federated learning may therefore be combined with secure aggregation, differential privacy, and retention controls in a defense-in-depth design whose guarantees depend on the complete protocol.

27 | Data lineage systems:

Apache Atlas and DataHub can capture metadata about data flows from pipeline execution, creating graphs in which nodes are datasets and edges are transformations. GDPR Article 30 requires records of processing activities (European Parliament and Council of the European Union 2016), not automated model-level lineage specifically. Lineage can nevertheless help identify candidate derived artifacts for deletion, restriction, retraining, or legal review when an applicable request arrives.



Retained audit events turn accountability into a growing storage workload.

of the European Union 2016). Compliance requirements transform legal obligations into system architecture constraints that shape pipeline design, storage choices, and operational procedures. GDPR's data minimization principle requires limiting collection and retention to what is necessary for stated purposes. Article 15 access rights cover personal data undergoing processing and specified information about that processing, not every artifact merely associated with a user.

Voice assistants operating globally face overlapping regulatory regimes because requirements vary by jurisdiction and apply differently based on user age and data sensitivity. GDPR cross-border-transfer rules permit transfers through adequacy decisions, appropriate safeguards, or limited derogations rather than imposing categorical data localization. These rules can still drive regional storage, replication, and processing choices. Data cards (Pushkarna et al. 2022) record provenance, intended uses, and risks as operational metadata, as illustrated in figure 15.6, but a valid card does not by itself make a dataset or model compliant.

15.5.4 Building data lineage infrastructure

Data-card fields become operational checks once they enter the pipeline: provenance, intended use, risk, and retention metadata determine which datasets can train which models and which artifacts must be traced during an audit. Compliance obligations are only as credible as the infrastructure that demonstrates them. When a regulator asks “which training data produced this model?” or a user invokes an applicable right to erasure, the organization must answer with engineering precision, not manual investigation. Data lineage provides this capability, transforming compliance documentation into operational infrastructure that powers governance across the ML lifecycle. Modern lineage systems like Apache Atlas and DataHub²⁷ can integrate with pipeline orchestrators such as Airflow and Kubeflow to capture relationships automatically. When an Airflow directed acyclic graph reads audio files from S3 and transforms them into spectrograms, the lineage system records each step and traces a feature back to its source audio file. Automated tracking also supports deletion requests. When a user invokes applicable erasure rights, the lineage graph identifies candidate derived artifacts—features, embeddings, and trained model versions—for the appropriate deletion, invalidation, restriction, unlearning, retraining, or legal-review workflow.

A production KWS system can implement lineage tracking across the data engineering lifecycle. Source audio ingestion can create lineage records linking each audio file to its acquisition method, supporting verification of consent requirements. Processing pipeline execution extends lineage graphs as audio becomes features and embeddings, and each transformation can record code versions and hyperparameters. Training jobs create lineage edges from feature collections to model artifacts, recording which data versions trained which model versions. When a voice assistant device downloads a model update, lineage tracking can record the deployment, supporting recall if training data is later discovered to have quality or compliance issues. Lineage captures provenance, but accountable operation also requires access history: who touched the data, when, and under which authority.

15.5.5 Audit infrastructure and accountability

Lineage tracks *what* data exists and *how* it transforms through the pipeline. Governance also requires knowing *who* accessed data and *when*, an accountability dimension lineage alone cannot provide. Audit systems record these access events, creating accountability trails for obligations such as HIPAA and the Sarbanes-Oxley Act (SOX) where applicable. Production ML systems can generate enormous audit volumes, requiring infrastructure such as tamper-evident, access-controlled logs, efficient indexing that supports queries for specific user or dataset accesses, automated analysis that detects anomalous patterns, and retention/deletion rules matched to the governing obligation.

KWS systems implement multi-tier audit architectures that balance granularity against performance and cost:

- **Edge devices:** Log critical events locally, with logs periodically uploaded to centralized storage for compliance retention.
- **Feature stores:** Log policy-relevant access with purpose-limited metadata, such as the requesting service, authorized subject identifiers, and feature groups retrieved.

Open Images Extended - More Inclusively Annotated People (MIAP) Dataset Download  Related Publication  This dataset was created for fairness research and fairness evaluations in person detection. This dataset contains 100,000 images sampled from Open Images V6 with additional annotations added. Annotations include the image coordinates of bounding boxes for each visible person. Each box is annotated with attributes for perceived gender presentation and age range presentation. It can be used in conjunction with Open Images V6.		
Authorship PUBLISHER(S) Google LLC INDUSTRY TYPE Corporate - Tech DATASET AUTHORS Candice Schumann, Google, 2021 Susanna Ricco, Google, 2021 Utsav Prabhu, Google, 2021 Vittorio Ferrari, Google, 2021 Caroline Pantofaru, Google, 2021 PUBLISHER(S) Google LLC FUNDING TYPE Private Funding DATASET CONTACT open-images-extended@google.com		
Motivations DATASET PURPOSE(S) Research Purposes Machine Learning Training, testing, and validation KEY APPLICATION(S) Machine Learning Object Recognition Machine Learning Fairness PROBLEM SPACE This dataset was created for fairness research and fairness evaluation with respect to person detection. See accompanying article  PRIMARY MOTIVATION(S) - Provide more complete ground-truth for bounding boxes around people. - Provide a standard fairness evaluation set for the broader fairness community. INTENDED AND/OR SUITABLE USE CASE(S) ML Model Evaluation for: person detection, fairness evaluation ML Model Training for: person detection, Object detection Also: Person detection: Without specifying gender or age presentations Fairness evaluations: Over gender and age presentations Fairness research: Without building gender presentation or age classifiers		
Use of Dataset SAFETY OF USE Conditional Use There are some known unsafe applications. UNSAFE APPLICATION(S)  Gender classification Age classification UNSAFE USE CASE(S) This dataset should not be used to create gender or age classifiers. The intention of perceived gender and age labels is to capture gender and age presentation as assessed by a third party based on visual cues alone, rather than an individual's self-identified gender or actual age. CONJUNCTIONAL USE Safe to use with other datasets KNOWN CONJUNCTIONAL DATASET(S) The data in this dataset can be combined with Open Images V6 KNOWN CONJUNCTIONAL USES Analyzing bounding box annotations not annotated under the Open Images V6 procedure.		
METHOD Object Detection METHOD Fairness Evaluation SUMMARY A person object detector can be trained using the Object Detection API in TensorFlow. SUMMARY Fairness evaluations can be run over the splits of gender presentation and age presentation. KNOWN CAVEATS If this dataset is used in conjunction with the original Open Images dataset, negative examples of people should only be pulled from images with an explicit negative person image level label. The dataset does not contain any examples not annotated as containing at least one person by the original Open Images annotation procedure. KNOWN CAVEATS There still exists a gender presentation skew towards unknown and predominantly masculine, as well as an age presentation range skew towards middle.		

Figure 15.6: Data Governance Documentation: The populated Open Images Extended—More Inclusively Annotated People card organizes authorship, motivation, intended and unsafe uses, conjunctional use, and method caveats in one artifact. Recording permitted and prohibited uses makes the card reviewable governance evidence, though it does not establish compliance by itself. Adapted from the Open Images Extended—MIAP Data Card in Pushkarna et al. (2022).

- **Training infrastructure:** Logs which jobs read which data partitions, supporting investigation and retention verification. Access logs alone cannot prove that deleted user data is absent from later model weights.

The tiers split audit work according to where evidence is generated, while preserving enough context to reconstruct access and deletion claims.

Regulatory requirements extend audit responsibility to prediction-time behavior. Answering “why was this specific applicant denied a loan?” requires enough context to reconstruct the decision, but not necessarily a permanent copy of every raw feature. Without inference-time logging, an audit trail may answer access questions without reconstructing which model, threshold, and evidence produced a specific decision. Production audit infrastructure should retain the minimum purpose-limited evidence required—for example, model version, decision threshold, output, reason codes,

28 | Audit trail: Audit integrity requires controls against unauthorized alteration, but retention and deletion duties vary by record type and law; append-only object storage with retention controls, or cryptographic hash chains, can provide tamper evidence without creating a universal rule that records are never deleted. A large platform may log billions of events daily; HIPAA's Security Rule also imposes multi-year documentation-retention obligations for required policies and procedures (U.S. Department of Health and Human Services 2005). Retention planning must distinguish records that must be preserved from personal data that must be minimized or deleted.

and selected input provenance—with access controls, retention limits, and identifiers appropriate to the governing rules. The goal is to make per-decision reconstruction a controlled query rather than a manual investigation.

Together, the four governance domains—security, privacy, compliance, and audit²⁸—form an enforcement layer that supports the other practices in this chapter. Data governance helps capture measurements, record actions, and preserve evidence for review. Without this infrastructure, responsible engineering remains aspirational; with it, teams can demonstrate specific controls and investigate whether commitments were met.

15.6 Fallacies and Pitfalls

Teams can still fail after assembling assessment frameworks, fairness metrics, explainability mechanisms, efficiency analyses, and governance infrastructure. A team may retrofit fairness after benchmark success, trust aggregate accuracy, or treat compliance evidence as paperwork rather than system behavior, drawing on intuitions from traditional software engineering where bugs are local and testing is deterministic. Recognizing these failure patterns early, before a fallacy shapes a design decision, is far cheaper than discovering it after deployment.

Fallacy: *Responsibility can be addressed after the system achieves technical objectives.*

Teams assume fairness constraints can be retrofitted once models demonstrate strong benchmark performance. In production, early architectural decisions constrain what interventions remain feasible. Amazon's recruiting tool (see section 15.2.1) illustrates this trap: attempted remediation did not provide confidence that the system would avoid discriminatory recommendations, and Amazon abandoned the project (Dastin 2018). Organizations deferring responsibility may face redesign, deployment with documented risks, or cancellation. Integrating fairness constraints early can be less costly than retrofitting them after data contracts, monitoring, and release gates are fixed.

Pitfall: *Relying on aggregate metrics to assess fairness.*

Engineers assume high overall accuracy indicates the system works well for all users. The Flaw of Averages (section 15.3.3) reveals this intuition fails: aggregate metrics can conceal large subgroup disparities (section 15.2.4). The loan approval analysis in section 15.3.3.1 showed a 30 percentage-point TPR gap, with qualified minority applicants rejected at 4× the majority-group rate. These disparities can persist undetected when standard monitoring tracks only aggregates. Production systems require disaggregated evaluation with application-specific thresholds, such as the 1.25× error-rate ratio or 5 percentage point TPR difference used in this example.

Fallacy: *Removing sensitive attributes from training data eliminates bias.*

Teams remove gender, race, and protected attributes expecting this ensures fairness. Other features can act as proxy variables when they correlate with sensitive characteristics. ZIP codes, purchase patterns, browsing history, college names, and language choices can carry indirect demographic signal. Amazon's system (see section 15.2.1) penalized terms such as “women's” and graduates of two all-women's colleges (Dastin 2018). A population-health study found that correcting a cost-based proxy would increase the share of Black patients receiving additional help from 17.7 percent to 46.5 percent (Obermeyer et al. 2019). Removing protected attributes does not by itself establish fairness.

Pitfall: *Treating documentation as sufficient accountability.*

Teams invest effort in model cards, then consider responsibility requirements satisfied. Documentation provides transparency (section 15.3.2) but not enforcement. A model card specifying “not validated for high-stakes decisions” has no effect when the system is repurposed for loan approvals without technical restrictions. Accountability requires operational integration: monitoring dashboards, documented subgroup-disparity alert thresholds, incident response procedures, and access controls preventing deployment beyond validated use cases.

Fallacy: *Responsible AI is primarily a legal compliance issue.*

Teams treat responsibility as external oversight rather than engineering practice. Engineering decisions made months before legal review constrain the solution space more than any compliance assessment. Architecture selection determines what fairness interventions are feasible, while data

pipeline design establishes whether disaggregated evaluation is even possible. As section 15.2.5 establishes, systems designed with responsibility as an engineering objective enable efficient validation; systems where responsibility is added at late-stage review face redesign or deployment with documented risks.

Pitfall: *Measuring the environmental impact of training but not inference.*

Public discourse focuses on training-run carbon, and engineers often follow this framing when assessing environmental responsibility. The illustrative TCO analysis in section 15.4.3 shows why this focus is incomplete: under its unbatched service-demand assumptions, the inference-to-training cost ratio is about 40:1. A model trained periodically but served millions of times daily can have a lifecycle footprint dominated by inference rather than training. For the recommendation system analyzed in table 15.16, training accounts for 1.9 percent of three-year costs while inference accounts for 73.5 percent. In this example, inference emits about 63× as much CO₂ as training. Engineers who optimize training efficiency while ignoring per-query inference demand can leave the larger term in this scenario unexamined. These values are scenario-dependent, but they show why lifecycle accounting must include both terms.

Fallacy: *Model weights are exempt from data governance and deletion requests.*

Teams often assume that once training data has been compiled into model weights, the data is gone and compliance obligations cannot reach the artifact. Models can memorize training data, and membership inference attacks may reveal whether a record appeared in training (Shokri et al. 2017). Whether particular weights constitute personal data or require action after an erasure request is case-specific. Engineering teams should track data-to-model lineage and evaluate deletion, restriction, retraining, machine unlearning, or compensating controls with legal and privacy specialists when a request may affect deployed artifacts (Cao and Yang 2015; Bourtole et al. 2021).

Across these failures, the recurring mistake is to treat responsibility as a document, metric, or late-stage review rather than a system property. Measurable constraints, continuous monitoring, and enforceable governance carry responsibility through the full lifecycle.

15.7 Summary

Responsible engineering is ML systems engineering done completely, not a separate discipline. The chapter traced a path from failure diagnosis through prevention to enforcement, beginning with the responsibility gap (the distance between technical performance and responsible outcomes) and demonstrating how proxy variables, feedback loops, and distribution shift can harm users while conventional metrics remain acceptable. The engineering response includes checklists that systematize predeployment assessment, fairness metrics that make disparities measurable, explanation mechanisms selected for applicable stakeholder and regulatory requirements, and monitoring infrastructure that can surface silent failures and connect signals to response.

Translating responsibility concerns into measurable properties makes them tractable. A justified, application-specific bound on disparity is testable; “be fair” is not. This translation extends beyond fairness: in the chapter’s illustrative TCO scenario, a 20 percent accelerator-service-demand reduction saves \$304K and is modeled to avoid 19 t of CO₂ under the stated assumptions. Documentation becomes model cards with explicit intended use and known limitations. Governance becomes access control, lineage tracking, and audit infrastructure that makes compliance evidence available rather than merely aspirational. Across these cases, abstract ethical obligations become concrete engineering requirements that can be specified, tested, monitored, and enforced.

The responsible engineering practices developed in this chapter are integral components of complete engineering, not external constraints layered onto technical work. Systems that ignore fairness, efficiency, transparency, or governance are technically incomplete. The same rigor applied to latency budgets and memory constraints must extend to fairness criteria, environmental impact, and applicable regulatory requirements. Integrating these considerations from system inception makes obligations testable, exposes trade-offs earlier, and makes failures easier to diagnose and govern in production.



Key Takeaways: Reliable for whom?

- **Aggregate correctness hides harm:** A model can look accurate in aggregate while one subgroup's error rate is 43.1× another subgroup's rate, as in the Face++ Gender Shades audit. Responsible evaluation therefore starts with disaggregated and intersectional slices, not aggregate accuracy alone.
- **Responsibility becomes testable through thresholds:** "Be fair" is not testable, but bounded disparity, documented intended use, and explainability requirements are. Translating values into measurable constraints lets teams review trade-offs among fairness, accuracy, latency, and cost.
- **Efficiency is a social constraint:** An efficient model can reduce energy and cost and broaden who can deploy it. In the chapter's illustrative scenario, the inference-to-training cost ratio is 40:1, making per-query optimization responsible engineering.
- **Monitoring must watch outcomes:** Bias and privacy failures can continue with green uptime dashboards because harmful predictions look operationally normal. Production monitoring must track subgroup outcomes, data lineage, feedback loops, and incident paths with the same rigor as latency regressions.
- **Governance has to be built in:** Model cards, datasheets, access controls, erasure workflows, human-review paths, and audit trails are technical infrastructure, not optional policy overlays. Regulations such as GDPR require capabilities that should be designed into the pipeline; later remediation may remain possible but can be narrower and costlier once the system is serving decisions.



A system that does exactly what it was told is dangerous precisely because the telling is never complete. Every objective a model is given is a specification with gaps, and an optimizer is a machine for finding them. It can reproduce the bias latent in its data, chase the proxy instead of the goal, and call the result success because nothing in the objective said otherwise. Responsible engineering is the discipline of writing back in what the specification left out. Constraints and monitoring reduce the risk that optimization amplifies harms encoded in the data. The constraint is the same kind the rest of the book imposed in latency and memory, except that here it protects people the objective never named, and a model that is fast and accurate while wrong about whom it serves has not failed less than one that crashes, only more quietly.




What's Next: From technique to philosophy

The chapter closes a circle that began with the iron law of ML systems (principle 3). The optimizations developed in chapter 10, chapter 11, and chapter 14 were motivated by performance, but they can also support responsibility. Efficiency can reduce energy and emissions, compression can lower deployment resource requirements, and outcome monitoring can surface bias. The engineering techniques overlap, but the responsibility lens changes which outcomes are measured. Chapter 16 connects these pieces into a coherent philosophy of engineering excellence.

SYNTHESIS



Conclusion




- 16.1 Synthesizing ML Systems
- 16.2 Thirteen Quantitative Principles
- 16.3 Principles in Practice
- 16.4 Future Directions
- 16.5 Journey Forward
- 16.6 Fallacies and Pitfalls
- 16.7 Summary

Purpose

What does mastering the full stack enable that expertise in any single layer cannot?

A single production decision can travel through the entire stack. A data pipeline decides which events count as training signal; that signal shapes the architecture that can learn the task; the architecture determines memory footprint and arithmetic intensity; those properties constrain hardware choice, quantization strategy, serving latency, drift monitoring, and governance obligations. Mastered individually, each layer is a valuable skill. Mastered together, they provide the ability to reason across boundaries. An engineer who understands only compression can shrink a model, but cannot predict whether the accuracy loss matters for the deployment context. An engineer who understands only serving can optimize latency, but cannot trace a performance regression to a data pipeline change three stages upstream. The discipline of ML systems engineering is the discipline of seeing these connections, where one team's optimization becomes another team's constraint. The principles governing these interactions, including constraint propagation, the memory wall, the training-serving inversion, dispatch overhead, communication cost, and the recurring cost of operating models in production, are not tied to any specific framework, hardware generation, or model family. Technologies will change, but the underlying physical constraints and trade-offs will endure. In D·A·M terms, what endures is the ability to look at a system that does not yet exist and reason about how its data, algorithm, and machine constraints will interact, where its bottlenecks will emerge, and which design decisions will prove irreversible. That D·A·M habit of thinking in systems rather than components is what separates an engineer who can build a part from one who can build the whole.





Learning Objectives

- Synthesize core ML systems principles into a framework for reasoning across Data, Algorithm, and Machine constraints
- Trace how data, architecture, compression, hardware, serving, operations, and governance decisions propagate constraints across an ML system
- Apply lighthouse-model reasoning to diagnose bottlenecks across cloud, mobile, edge, recommendation, and TinyML deployments
- Evaluate deployment trade-offs using latency budgets, memory movement, drift, responsibility, and sustainability constraints
- Design a systems engineering posture for emerging contexts before fleet-scale coordination costs dominate

16.1 Synthesizing ML Systems

IMAGINE deploying a new image classification model to a fleet of mobile devices. The architecture team chose depthwise separable convolutions for efficiency. The compression team quantized to INT8 for speed. The serving team hit a p99 latency target of 50 ms. Each team succeeded by its own metric, yet within weeks, user complaints arrive: accuracy has dropped by four percentage points on specific firmware and device cohorts. The cause is a subtle interaction between the quantization scheme and a firmware-specific image preprocessing path. No component is broken in isolation, but the data pipeline, architecture, compression strategy, hardware target, and monitoring infrastructure have coupled in production.

Responsible engineering is not an external layer added after optimization, but the discipline of specifying, testing, monitoring, and governing the whole system. That lesson now generalizes across this foundational material. ML systems are a different engineering problem from traditional software because the model is inseparable from the system that produces, serves, and monitors it.

The book began with the iron law of ML systems (principle 3). Its three terms—data movement, compute, and overhead—which once seemed abstract, now serve as primary engineering levers for quantitative analysis of systems that once seemed opaque. Building intelligence requires both algorithms and adherence to the silicon contract (principle 4), the physical and economic agreement between the model and the machine. Arithmetic intensity and roofline reasoning convert vague performance intuitions into quantitative engineering decisions (Williams et al. 2009).

The quantitative foundation leads to a broader point: contemporary artificial intelligence achievements are an emergent property of D·A·M co-design, not any single algorithmic insight. Machine learning belongs to the same engineering tradition that built reliable computers, where emergent capabilities arise from coordinating many parts together. The transformer architecture introduced an attention-based model family (Vaswani et al. 2017), and later large language model systems such as GPT-3 and Llama 2 show how that family scaled into a central workload for modern ML systems (Brown et al. 2020; Touvron, Martin, et al. 2023). Its mathematical design alone does not explain its practical utility. That utility depends on integrating attention mechanisms with distributed training infrastructure, memory-efficient optimization techniques, and reliable operational frameworks.

Integration has concrete consequences. We often speak of the “model” as the weights file, a 500 MB blob of floating-point numbers. In a production environment, however, the weights are only one component of the true model, and often not the most important one. A model that produces perfect predictions is useless if it receives corrupted inputs, and a model that trains flawlessly will fail if it cannot be deployed reliably. The *true model* is the sum of the data pipeline that defines what the model sees, the training infrastructure that determines what it learns, the serving system that decides how it interacts with the world, and the monitoring loop that keeps it tethered to reality. Optimize the system, and the model improves. Neglect the system, and the model degrades. Systems engineering is not a wrapper around ML; it is the implementation of ML. *The system is the model.*

📌 Checkpoint 16.1: Systems thinking

The system's dependencies, feedback loops, and request path make its boundaries testable.

The integration

- ❑ **Dependencies:** can you trace a data-pipeline change through training and serving to identify its latency effect?
- ❑ **Feedback loops:** have you mapped how the model's predictions influence its future training data?

The holism

- ❑ **End-to-end:** can you trace a user request from the UI, through the network, preprocessing, model, postprocessing, and back to the UI?

Tracing a request end to end, as the checkpoint asks, makes the same point structurally: system boundaries define model capabilities. That insight has guided the exploration throughout this book. The arc began by making the substrate explicit: data engineering (chapter 4) and data selection (chapter 9) determined what the system could learn, neural computation (chapter 5) and network architectures (chapter 6) determined how that signal became computation, and training systems (chapter 8) and frameworks (chapter 7) turned the computation into an executable optimization process.

Once the substrate existed, the engineering problem shifted from building to renegotiating constraints. Model compression (chapter 10) changed the accuracy-memory-latency trade-off; hardware acceleration (chapter 11) tested whether the resulting computation could actually feed the silicon; benchmarking (chapter 12) supplied the measurement discipline needed to distinguish real speedup from artifact. Production then exposed the assumptions that survived the lab but failed under load: serving systems (chapter 13) had to meet latency budgets, operational practices (chapter 14) had to keep models healthy as distributions shifted, and responsible engineering (chapter 15) had to evaluate whom the system served and where aggregate performance concealed subgroup failure. These chapters fill out the introduction's five-pillar framework: data engineering, training systems, deployment infrastructure, operations and monitoring, and the ethics-and-governance pillar that threads Part IV rather than standing apart from it.

Each chapter contributed a piece. The real lesson, however, lies not in any individual piece but in *how* the pieces constrain each other. An architecture choice enabled a compression choice, which enabled an acceleration choice, which shaped a serving constraint, which defined an operational requirement. MobileNetV2's depthwise-separable design targeted efficient mobile vision inference (Sandler et al. 2018), while integer-arithmetic quantization made INT8 deployment a practical inference path (Jacob et al. 2018). That combination can enable mobile NPU deployment, shape a p99 latency constraint, and require drift monitoring across heterogeneous device populations. Every decision propagated forward, and the engineer who understands only one layer cannot predict how changes ripple through the rest.

The lighthouse models provide a constraint map for reasoning about ML systems as wholes rather than as collections of parts. They trace the same interactions across chapters before the synthesis formalizes thirteen quantitative principles, including exact bounds and assumption-dependent diagnostic models, for reasoning about ML system behavior. Those principles then carry into three application domains, future directions where systems thinking will matter most, and the engineering responsibility that accompanies building systems of this power.

Lighthouse models: Constraint propagation

The five lighthouse models introduced in section 1.7 made this constraint propagation concrete, serving as systems detectives throughout the book. Each revealed how different workloads expose different bottlenecks.



Architecture choices cascade into compression, serving, drift, and governance.

The five lighthouse workloads expose distinct constraint regimes:

- **ResNet-50:** Batch size can turn image inference from a memory-bound path into compute-bound throughput.
- **GPT-2/Llama:** Low-batch autoregressive decode is often bandwidth bound because each step reuses weights too little; batching, prefill, architecture, parallelism, and hardware can shift the bottleneck.
- **MobileNetV2:** Depthwise separable convolutions and INT8 quantization trade representational capacity for mobile NPU deployment in a power-constrained regime.
- **DLRM:** Terabyte-scale embedding tables can make memory *capacity* a binding constraint alongside memory *bandwidth*, forcing engineers to design around where data physically resides and how sparse operations behave.
- **Keyword spotting (KWS)/Wake Vision:** Sub-megabyte microcontroller models with always-on inference under milliwatt power budgets make every byte and milliwatt matter.

Together, these five workloads span the deployment spectrum from data center to microcontroller, probing the bottlenecks these principles diagnose and testing the optimization strategies developed throughout the book. The systems thinking we developed by tracing these lighthouses across chapters, from architecture design through training, optimization, and deployment, is the integrated perspective that distinguishes ML systems engineering from isolated algorithm development.

Table 16.1 traces this journey for a single model, MobileNetV2, showing how principles from across the book converge on one engineering artifact. The table walks through seven phases (from foundational constraints through architecture, training, compression, acceleration, serving, and operations) showing how each phase's decisions propagate forward to shape what becomes possible in subsequent phases.

Table 16.1: The Lighthouse Journey (MobileNetV2): Tracing one model through seven phases of the systems stack shows how decisions in one domain (for example, architecture) propagate constraints and opportunities into downstream domains (for example, hardware acceleration and monitoring).

Journey Phase	System Lens	MobileNetV2 Implementation
Foundations (chapter 1)	The AI Triad	Bounded by machine constraints (Battery/Thermal)
Architecture (chapter 6)	Algorithmic Efficiency	Depthwise Separable Convolutions: 8.7× fewer FLOPs for a representative 3-by-3, 256-output-channel layer and 13.7× fewer operations than ResNet-50 at ImageNet scale
Training (chapter 8)	Throughput vs. Latency	Optimized for single-request mobile latency; data augmentation can improve robustness
Compression (chapter 10)	Navigating the Pareto Frontier	INT8 Quantization: FP32 uses 4× and FP16 uses 2× the per-value storage of INT8, with accuracy revalidated per deployment
Acceleration (chapter 11)	Honoring the Silicon Contract	Mapping kernels to Mobile NPUs (for example, Apple Neural Engine) to maximize hardware utilization
Serving (chapter 13)	Respecting the Latency Budget	p99 < 50 ms constraint; optimizing preprocessing (resize/normalize) to avoid CPU bottlenecks
Operations (chapter 14)	Managing System Entropy	Drift Monitoring: Detecting distribution shifts and tracking cohort-level accuracy across heterogeneous device populations and lighting conditions

The table shows how decisions in one row constrain the options in the next. Architecture choices (depthwise separable convolutions) enabled compression choices (INT8 quantization), which in turn enabled acceleration choices (mobile NPU deployment). Constraint propagation recurs across ML systems, and the MobileNetV2 journey is one instance of that structure. The question is which quantitative tools recur across specific models and technologies. The answer lies in thirteen quantitative principles, some exact bounds and others assumption-dependent diagnostics.

16.2 Thirteen Quantitative Principles

Throughout this book, each part introduced quantitative tools for reasoning about ML system behavior. These **thirteen quantitative principles** deliberately mix exact mathematical bounds

with engineering decompositions, fitted local models, policy requirements, and design heuristics. Table 16.2 collects all thirteen in one place, organized by the four parts that revealed them. Their value comes from applying each within its stated assumptions, not from treating every row as a universal law.

Table 16.2: Thirteen Quantitative Principles: Each tool appears in the part where its governing constraint or diagnostic use first becomes visible. The collection combines bounds, decompositions, fitted models, policy requirements, and heuristics. Each row is useful only under its stated assumptions; together they provide a shared analytical vocabulary for system design, optimization, and deployment rather than a set of universal invariants.

#	Principle	Part	Core Equation/Statement	What It Predicts
1	Data-as-Code Principle	I: Foundations	Behavior = $f(\text{data, algorithm, code, randomness})$	Data changes behavior; other inputs also matter
2	Data-Gravity Principle	I: Foundations	Move compute toward data when repeated transfer costs exceed placement costs	Depends on volume, reuse, network, and compute mobility
3	Iron Law of ML Systems	II: Build	$T_{\text{seq}} = D_{\text{vol}}/\text{BW} + O/(R_{\text{peak}}\eta_{\text{hw}}) + L_{\text{lat}}$; overlap ranges from max to sum	Stages may add or overlap
4	Silicon Contract	II: Build	$I_{\text{ridge}} = R_{\text{peak}}/\text{BW}$; compare I_{model} with I_{ridge}	Diagnoses bandwidth- versus compute-limited operation
5	Pareto Frontier	III: Optimize	$\nexists c' \neq c : [\forall k M_k(c') \geq M_k(c)] \wedge [\exists j M_j(c') > M_j(c)]$	No distinct configuration dominates a frontier point
6	Arithmetic Intensity Law	III: Optimize	$R_{\text{attain}} \leq \min(R_{\text{peak}}, I \times \text{BW})$	More compute cannot raise a bandwidth ceiling
7	Energy-Movement Invariant	III: Optimize	$E_{\text{total}} = \sum_j N_j E_j$; DRAM/FLOP cost ratio: 173–582×	Total energy depends on event counts and costs
8	Amdahl’s Law	III: Optimize	$\text{Speedup} = \frac{1}{(1-f_{\text{parallel}}) + \frac{f_{\text{parallel}}}{S_{\text{parallel}}}}$	The serial fraction caps all parallelism gains
9	Verification Gap	IV: Deploy	$\Pr_{(X,Y) \sim P_{\text{deploy}}}[d(f(X), Y) \leq \tau] \geq 1 - \epsilon$	Specify distance, tolerance, population, and confidence
10	Statistical Drift Diagnostic	IV: Deploy	$\text{Accuracy}(t) \approx \text{Accuracy}_0 - \lambda \mathcal{D}(P_t \ P_0)$	A local fit; drift need not lower quality
11	Training-Serving Skew Diagnostic	IV: Deploy	$S_{\text{skew}} = \mathbb{E}_{X \sim P_{\text{deploy}}}[d(f_{\text{serve}}(X), f_{\text{train}}(X))]$	Output mismatch signals risk, not accuracy loss
12	Latency Budget Principle	IV: Deploy	$T_q \leq L_{\text{budget}}$	The product SLO selects q and its budget
13	Bias Feedback Model	IV: Deploy	$\Delta_g(k) \approx \Delta_g(0)\alpha_{\text{fb}}^k$ with fitted α_{fb}	Feedback may amplify harm; measure and intervene

The thirteen principles are not independent axioms. They form an integrated framework connected by a single meta-principle: the conservation-of-complexity heuristic.¹ Complexity removed from one interface often reappears in another part of the system, but this is not a physical conservation law and does not imply that every simplification has an equal compensating cost. Its value is diagnostic: after simplifying one component, check where validation, state, coordination, or operational burden changed. The test is whether the principles explain the same lighthouse bottlenecks from data, model, hardware, and deployment perspectives without contradicting one another.

Foundations: Where complexity originates (principles 1–2)

The data-as-code principle (1) and the data-gravity heuristic (principle 2), established in Part I and developed in chapter 4, establish data as a major logical input and potential physical anchor. Behavior also depends on algorithm and implementation, while compute-to-data placement depends on volume, reuse, network cost, and compute mobility. Model behavior and architecture therefore inherit constraints from the data substrate.

The lighthouse models illustrate both principles directly. ResNet-50 and GPT-2 depend on both their architectures and their training data. DLRM’s terabyte-scale embedding tables can make a strong case for designing the system around where the data physically resides. These principles help explain why the compute-to-data pattern recurs across deployment contexts without turning it into a universal placement rule.

¹ | **Conservation-of-complexity heuristic:** Tesler’s Law, a design aphorism, says that an application’s irreducible complexity must be handled somewhere in the interaction among user, application, and platform (Tesler 1984); extending it to all ML-system complexity is an analogy, not a physical law. Quantization may add validation burden, and abstraction may move implementation detail behind an interface, but good design can also remove accidental complexity outright. Large language model application pipelines illustrate a possible shift: simplifying the user-facing interface with shorter or vaguer prompts may move work into system prompts, retrieval, or output verification. Use the heuristic to search for displaced costs, not to assume that an equal compensating cost must exist.

Build: How complexity becomes computation (principles 3–4)

The iron law (principle 3) and the silicon contract (principle 4) guide decisions in constructing an ML system. The iron law's three-term decomposition (introduced in section 1.7) identifies which lever to pull; the silicon contract determines which term dominates for a given architecture-hardware pair. As the lighthouse journey showed, each model can expose a different constraint regime: batched ResNet-50 inference can be compute bound, low-batch Llama decode can be bandwidth bound, DLRM can be capacity bound, and MobileNetV2 reshapes its computation to fit mobile NPU constraints. Section A.5.1 maps each of these regimes to the optimizations that pay off and the ones that waste effort, turning the diagnosis of compute-bound versus bandwidth-bound versus capacity-bound into an action plan. Chapter 8 confirmed that training time falls most when engineers optimize the dominant term rather than distributing effort uniformly.

Optimize: How constraints shape trade-offs (principles 5–8)

The four optimization principles form a tightly coupled diagnostic chain. The Pareto frontier (principle 5) identifies nondominated trade-offs after objective directions are normalized: quantization trades precision for memory traffic, pruning trades capacity for speed, and distillation trades training compute for inference efficiency. The arithmetic intensity law (principle 6) diagnoses whether compute or bandwidth sets the ideal ceiling. The energy-movement invariant (principle 7) combines per-event costs with event counts: in the book's reference constants, one DRAM access costs about $173\text{--}582\times$ as much as one FP32/FP16 arithmetic operation, but workload-total dominance depends on how many of each occur. Amdahl's Law (principle 8) sets the ceiling on any parallelism gain, explaining why data loading and preprocessing can become bottlenecks in highly optimized systems.

MobileNetV2 (our lighthouse from chapter 6) navigates all four simultaneously: depthwise separable convolutions reshape the Pareto frontier (Sandler et al. 2018), INT8 quantization exploits the arithmetic intensity law by increasing operations per byte through reduced memory traffic (Jacob et al. 2018), and the resulting energy savings respect the energy-movement invariant while Amdahl's Law explains why a preprocessing stage that remains unaccelerated can limit end-to-end speedup. The KWS lighthouse pushes these trade-offs to their extreme, where sub-megabyte models on microcontrollers leave zero margin for waste on any axis.

Deploy: How reality defeats assumptions (principles 9–13)

The deployment principles address failures that bench testing cannot rule out: a system can work correctly on the bench yet fail silently in production. The verification gap (principle 9) means statistical verification holds only over a stated deployment population, task distance, tolerance, target error, and finite-sample confidence procedure; testing estimates behavior rather than proving correctness for every future input. The statistical drift diagnostic (principle 10) detects distribution change but does not determine whether accuracy worsens, stays constant, or improves, so a local degradation curve is valid only when fitted against outcomes. The training-serving skew diagnostic (principle 11) is likewise a risk indicator rather than a universal accuracy-loss equation, although preprocessing or numerical differences can still cause quality changes. The latency budget (principle 12) constrains serving at the quantile selected by the product requirement, which may be p95, p99, or another tail measure. Finally, the bias feedback model (principle 13) can represent disparity amplification, but an exponential recurrence requires a measured, approximately constant feedback factor and no effective intervention.

The five deployment principles explain why chapter 14 devoted extensive attention to monitoring, drift detection, feature stores, and disaggregated subgroup metrics: operational infrastructure must surface silent failures and connect evidence to response. A DLRM recommendation system that achieves excellent offline accuracy still needs parity checks when training-serving skew corrupts feature values (principle 11) and outcome checks when user behavior drifts seasonally (principle 10). GPT-2/Llama serving must respect its selected latency quantile through techniques such as continuous batching and speculative decoding, as detailed in chapter 13, because excessive response time may violate the product requirement. A loan approval system can satisfy every other principle while still systematically denying credit to underserved communities, and feedback can compound the harm unless disaggregated subgroup monitoring detects it and triggers review.

The integrated framework

The thirteen principles are not a checklist to apply sequentially. They form a web of mutual constraints. The conservation-of-complexity heuristic prompts engineers to look for where a local simplification changes burdens elsewhere.

We can trace these mutual constraints concretely by following what happens when an engineer quantizes a model from FP16 to INT8. This single decision navigates the Pareto frontier (principle 5), trading precision for memory traffic. The consequences do not stop there: quantization changes the model's silicon contract (principle 4), shifting where it sits on the arithmetic intensity curve (principle 6) and altering its energy profile (principle 7). When that quantized model is deployed, the latency budget (principle 12) governs whether the speedup meets the service-level objective (SLO), while deployment validation must verify that the quantized serving path preserves the behavior accepted during compression testing. A single quantization decision ripples through the Pareto frontier, silicon contract, and latency budget simultaneously, where a win in one (memory traffic) must be validated against a risk in another (numerical error).

That trace does not require every principle to apply at once. It shows how the relevant principles become active as a decision moves from model representation to hardware execution to production validation. Data placement affects where the model can run, Amdahl's Law limits how much the faster kernel can improve the whole request path, verification bounds the resulting accuracy loss, and outcome monitoring tests whether the validated behavior persists after deployment. The engineer's task is to trace displaced costs rather than assume that complexity is conserved.

A small deployment proposal makes this web of constraints concrete.



Checkpoint 16.2: Applying the principles

A colleague proposes quantizing your model from FP32 to INT8 to reduce serving costs.

Trace the principles

- ☐ **Pareto frontier** (principle 5): What accuracy are you trading for the bandwidth gain?
- ☐ **Silicon contract** (principle 4): Does the target hardware efficiently accelerate the required INT8 operators?
- ☐ **Serving-path validation**: does the quantized serving artifact preserve the behavior accepted during compression testing?
- ☐ **Latency budget** (principle 12): Does the speedup bring you within SLO, or create head-room for batching?

To see this cycle of mutual constraint in action, trace the flow in figure 16.1. The four phases (Foundations, Build, Optimize, Deploy) surround a central hub representing the conservation-of-complexity heuristic, and the arrows map the flow of engineering decisions: each phase's choices constrain what becomes possible in the next, and the cycle eventually feeds back to the beginning. Decisions in Build constrain Optimize, while production evidence such as drift, skew, and outcome changes can feed back into Foundations. The engineer's role is to manage this flow, ensuring that displaced burdens land where they can be handled efficiently.

The Deploy-to-Foundations feedback arrow is central to this cycle. Principles nine through thirteen expose signals and constraints that may require a corrected release, fallback, new data, retraining, or a fresh optimization pass. When one appears, engineers must diagnose which response fits the cause rather than react automatically to a drift alarm. The cycle operates within the single-system scope of this book: the goal is not to name every future architecture, but to make feedback visible early enough that engineers can redesign before failures compound.

Tracing a quantization proposal through four principles is one diagnostic pass; the same habit applies when the bottleneck is not an optimization proposal but the cost of serving a single generated token.




Figure 16.1: The Cycle of ML Systems (13 Principles): A four-phase systems engineering lifecycle organized around the conservation-of-complexity heuristic. The phases connect in a feedback cycle, and the thirteen principles are grouped along the transitions as bounds, decompositions, fitted models, policy requirements, and design heuristics whose assumptions must be checked.

Napkin Math 16.1: The cost of a token

Problem: When we apply the iron law (principle 3) and the arithmetic intensity law (principle 6) to serving one token from a 70-billion-parameter model like Llama 2 70B, with its FP16 weights sharded evenly across two NVIDIA H100s, which idealized lower bound is larger: memory transfer or peak compute? Section D.3.2 supplies the H100 specifications. This calculation excludes KV-cache traffic, activations, context-dependent attention-over-KV FLOPs, interconnect communication, and dispatch overhead.

Physics:


- **Model-weight byte volume moved** (D_{vol}): 70 billion parameters $\times$ 2 bytes (FP16) = 140 GB.
- **Compute** (O): $O \approx 2 \times P = 140$ GFLOP per token, where P is the parameter count.
- **Hardware:** Two H100s with aggregate BW = 6.70 TB/s, $R_{\text{peak}} \approx 1978$ TFLOP/s FP16.

Math:

- **Time to move data:** $T_{\text{mem}} = \frac{140 \text{ GB}}{6700 \text{ GB/s}} \approx 20.9 \text{ ms}$
- **Time to compute:** $T_{\text{comp}} = \frac{140 \times 10^9 \text{ FLOP}}{1978 \times 10^{12} \text{ FLOP/s}} \approx 0.07 \text{ ms}$

Systems insight:

The idealized memory-transfer lower bound T_{mem} is 295.2 $\times$ larger than the peak-compute lower bound T_{comp} . Under this batch-one model, decode sits left of the roofline ridge in the bandwidth-bound regime. Batching can increase reuse, while quantization can reduce weight traffic. Optimizing only compute execution can reduce realized compute time toward the 0.07 ms lower bound, but it cannot reduce the 20.9 ms memory-transfer lower bound.



In this batch-one lower-bound model, decode sits left of the roofline ridge in the bandwidth-bound regime.

This calculation shows the framework operating as a diagnostic instrument rather than an abstract taxonomy. The chapters applied these bounds, models, and heuristics to specific engineering

decisions, often without naming them explicitly. Tracing those applications across three domains—building foundations, engineering for scale, and navigating production reality—reveals how the framework has guided the analysis throughout the book.

16.3 Principles in Practice

A team that memorizes all thirteen principles but cannot identify their assumptions or apply them to a real deployment decision has learned nothing. The test is the same across the three domains that span the ML lifecycle: building technical foundations, engineering for scale, and navigating production reality. Systems thinking connects what isolated component analysis cannot.

Building technical foundations

The data-as-code principle (1) shaped chapter 4, echoing Karpathy’s Software 2.0 framing of datasets and model architecture as source code (Karpathy 2017) while recognizing that algorithms and serving code also affect behavior. Mathematical foundations (chapter 5) established the computational patterns relevant to the silicon contract: the matrix multiplications at the heart of neural computation determine arithmetic intensity, whose position relative to the hardware ridge point helps diagnose whether a workload is bandwidth or compute limited. Framework selection (chapter 7) illustrated the silicon contract’s practical consequence: each framework constrains graph optimization, memory management, hardware backend support, and the deployment paths that remain open. An engineer who selects a framework without considering those implications may discover too late that the chosen path forecloses the most efficient deployment option.

Foundational choices (what data to curate, which computational primitives to rely on, which framework to adopt) propagate into later engineering decisions. That propagation becomes especially visible when a system must scale beyond a single machine, where the iron law’s three terms expand from chip-level quantities to cluster-level constraints.

Engineering for scale


Training systems (chapter 8) demonstrated the iron law in action: data parallelism can reduce per-step compute time by distributing work across GPUs while adding communication, mixed precision can reduce data movement for tensors represented in FP16 rather than FP32, and gradient checkpointing trades recomputation for memory capacity, each technique pulling a different lever of the same three-term equation. Model compression (chapter 10) navigated the Pareto frontier directly: MobileNetV2’s INT8 quantization and DLRM’s embedding pruning each traded one metric for another, while the arithmetic intensity law diagnosed which trade-off would yield the greatest return for a given hardware target.

Building and optimizing a model, however, is only half the engineering challenge. The other half begins the moment the model leaves the training cluster and enters production, where statistical requirements, fitted diagnostics, and SLO policies govern behavior and where the optimizations that worked on the bench must survive the unpredictability of real-world traffic.

Navigating production reality

The transition from *training* to *inference* often changes optimization objectives: where training maximizes throughput over days, interactive inference may optimize a selected tail-latency quantile in milliseconds. The illustrative scenario uses a 50 ms mean and a 2,000 ms p99, so the selected tail is 40× the mean; another product may specify p95, p99.9, or separate deadlines by request class. Tail-latency analysis shows why mean latency can hide the requests that dominate user experience (Dean and Barroso 2013). In ML serving, tail spikes can arise from garbage collection pauses in Python-based inference frameworks, dynamic batching that forms suboptimally under uneven traffic, or autoregressive generation in which a long prompt or output drives much more work than a typical short response.

Beyond technical performance, chapter 15 broadened the framework to include societal impact. Verification requirements motivate monitoring for fairness violations alongside performance: tracking prediction distributions across demographic groups, detecting bias amplification over time (principle 13), and alerting on unexplained accuracy disparities. Drift monitoring also applies to



Mean latency hides the tail; the selected tail quantile governs the SLO.

subgroup distributions and outcomes, where accuracy may change for underrepresented populations even as aggregate metrics remain stable. Responsible AI is therefore an integral dimension of systems engineering, a first-class design constraint governed by the same measurement discipline as performance.

The three domains demonstrate that the thirteen principles are working tools rather than universal axioms. The question now is *where* their assumptions will be tested next, as ML systems expand into new deployment contexts, confront new failure modes, and pursue increasingly ambitious goals.

16.4 Future Directions

The framework is most useful when it forecasts where constraints may bind next. Three areas put the same physics under increasing pressure: deployment across diverse contexts, robustness under adversarial conditions (Goodfellow et al. 2015), and societal applications whose failures carry public consequences. A fourth horizon, systems that compose multiple models, tools, and verifiers or grow beyond one machine, extends the same lens rather than replacing it.

Applying principles to emerging deployment contexts

Deployment diversity tests whether one quantitative framework can explain systems with contrasting resource regimes. The cloud offers abundant power and centralized hardware, edge and mobile devices operate under latency and battery budgets, and TinyML and embedded systems compress the same design problem into kilobytes and milliwatts. Generative AI is not a fifth deployment environment; it is a workload class that stresses all four.

In the cloud regime, the binding decision is how to turn abundant hardware into useful throughput without letting data movement, capacity, or cost dominate. Dense workloads such as ResNet-50 chase GPU utilization through kernel fusion, mixed precision training, and gradient compression, while DLRM-style recommendation systems must also manage embedding-table capacity, placement, and sparse access patterns. Chapter 10 and chapter 8 explored these techniques, demonstrating how they combine to balance performance optimization with cost efficiency at scale.

In contrast, mobile and edge systems face stringent power, memory, and latency constraints that demand sophisticated hardware-software co-design. Efficient architectures introduced in chapter 6 (such as depthwise separable convolutions and neural architecture search) combined with compression techniques from chapter 10 (such as quantization and pruning) enable deployment on devices where the book's reference mobile platform has about 10× smaller memory capacity and about 140× smaller power envelope than an H100-class accelerator. Edge deployment matters when latency, privacy, connectivity, energy, or per-request cost make centralized serving the wrong abstraction; in those regimes, efficiency becomes part of accessibility rather than a separate optimization.²

Autoregressive generative models, illustrated by the GPT-2/Llama lighthouse family, stress the same constraints at token-serving scale. Low-batch decode for dense autoregressive models is often bandwidth bound because each step reuses each weight too little; larger batches can raise arithmetic intensity, while prefill is commonly compute intensive. Model partitioning across devices (splitting one model across multiple accelerators), which extends the parallelism chapter 8 previewed, redistributes weights and work while adding communication. Speculative decoding (chapter 13) trades additional compute for lower decoding latency. Together, these mechanisms demonstrate how the principles adapt as workload structure changes.

At the opposite extreme, TinyML and embedded systems, the domain of our KWS/Wake Vision lighthouse, face kilobyte memory budgets, milliwatt power envelopes, and potentially long deployment lifecycles. Success in these contexts depends on the full systems engineering approach: careful measurement reveals actual bottlenecks, hardware co-design improves efficiency, and planning for failure supports reliability despite severe resource limitations. Resource constraints have driven efficient architecture families such as MobileNets (Howard et al. 2017; Sandler et al. 2018) and EfficientNets (Tan and Le 2019) that also inform broader model-efficiency practice, demonstrating how systems constraints can catalyze algorithmic innovation.

The same physical bounds apply across these paradigms, while fitted statistical models and SLO policies remain deployment-specific. Success depends on checking those distinctions and applying the principles together rather than pursuing isolated optimizations. The more deployment contexts a

² | **AI democratization:**

Making AI accessible beyond a small number of well-resourced organizations through efficient systems engineering. Mobile-optimized models and cloud APIs can widen access, but doing so sustainably requires systematic optimization across hardware, algorithms, and infrastructure to maintain quality at scale.

system spans, the more failure surfaces it creates. Robustness, not coverage alone, therefore becomes a binding constraint at the next frontier.

Building robust AI systems

ML systems can respond confidently and incorrectly while ordinary availability checks stay green, and no one may notice for weeks. Distribution shifts may alter accuracy without code changes, adversarial inputs can exploit vulnerabilities invisible to standard testing, and edge cases can reveal training-data limitations that debugging alone cannot fix. These are credible production risks, not outcomes guaranteed by a single divergence metric.

Finite testing estimates behavior on a defined population with uncertainty; it cannot prove correctness for every future input. Drift signals show when that population may have changed, while labeled outcomes determine whether quality changed. Together they motivate continuous monitoring as a design requirement, along with fallback policies and periodic revalidation, without claiming that every distribution shift causes degradation. The operational question is whether a failure will be detected and diagnosed before its impact spreads.

Robustness therefore demands designing for graceful degradation, not only prevention. At the single-system scale, that discipline appears as fallback paths, uncertainty thresholds, version-specific rollback policies, and monitoring hooks. Rollback addresses a bad release; external drift may instead call for alerting, traffic reduction, fallback, data collection, or retraining after outcome evidence confirms harm. At larger scale, the same logic extends to hardware redundancy and ensemble-style diversity. As AI systems assume increasingly autonomous roles in healthcare, transportation, and finance, the gap between “works in the lab” and “works in the world” becomes the critical engineering challenge. Robustness becomes more essential as systems add components, because each interface creates another timeout, stale input, inconsistent state, or recovery path to monitor.

AI for societal benefit

Robust systems are the prerequisite for deploying AI in domains where technical failures carry public consequences. A medical AI that fails unpredictably cannot be trusted with patient care. An educational system that degrades under load cannot serve the students who need it most. A climate model that produces confident but uncalibrated predictions may misdirect policy decisions affecting millions of lives. In each domain, the thirteen principles provide shared questions and bounds, while domain evidence determines acceptable policy.

Each domain stresses a different governing constraint before it can deliver social value:

- **Scientific discovery:** Protein folding, drug interaction modeling, and materials science can require substantial training or inference throughput governed by the iron law (principle 3) and silicon contract (principle 4); when work is distributed, coordination adds another systems cost.
- **Healthcare AI:** Clinical review, calibrated uncertainty, external validation, and continuous monitoring can become safety-critical because a diagnostic model trained on one hospital’s population may degrade when deployed to another with different demographics, disease prevalence, or imaging equipment.
- **Personalized education:** Protecting sensitive student-interaction data used for personalization at global scale stresses data governance and the data-as-code principle (1); interactive services must also meet their selected latency budgets.

All three applications demonstrate that technical excellence alone is insufficient. The principles developed throughout this book (the D-A-M taxonomy, the thirteen quantitative tools, and the integrated reasoning framework) provide the systems engineering *foundation*, but the *application* of that foundation requires domain knowledge that no single discipline can supply.

The bounded nature of these applications is what makes their systems constraints tractable: a diagnostic medical model classifies conditions within a defined label set, and a climate model projects climate statistics within physical constraints. The next frontier asks whether the same principles can guide systems that delegate work across multiple components while preserving end-to-end requirements.

System composition as a stress test

The most ambitious stress tests for these principles are systems whose task boundaries are not fixed in advance. A task-general assistant or multi-component ML service may route one request through retrieval, planning, tool execution, generation, and verification. The governing challenge is systems engineering as much as algorithm design: the surrounding system must bound latency, reliability, cost, safety, and observability as work fans out across components.

System composition makes several central principles active at once:

- **Iron law:** The computation that each component performs must be budgeted as work fans out across retrieval, planning, tool execution, generation, and verification.
- **Silicon contract:** The system must honor hardware-specific constraints across CPUs, NVIDIA H100-class GPUs, Tensor Processing Units, and custom accelerators.
- **Pareto frontier:** The trade-off surface expands from two or three metrics, such as accuracy, latency, and memory, to a larger surface that also includes safety, fairness, factuality, privacy, and cost.
- **Statistical drift:** Drift applies not only to the final output, but also to retrieved documents, tool responses, and intermediate decisions.

A composed system cannot rely on a single model-quality claim; it needs interfaces whose behavior can be measured, because each interface is where one component's assumptions about another become testable (Lampson 1983). Composed systems trade monolithic simplicity for explicit coordination. A retrieval component finds relevant information, a reasoning component processes it, a tool call may query an external system, and a verifier checks the output. Each step can be independently updated, monitored, and debugged, but each also creates another interface contract. The decomposition trades latency and architectural complexity for control and observability, an example of a Pareto trade-off and possible complexity shift rather than a conservation law.

The systems cost is visible in a single request. If an assistant fans out to retrieval, a planner, two tools, a generator, and a verifier, its realized latency follows the request's critical path, with parallel tool calls contributing their maximum rather than their sum, plus orchestration overhead. Its reliability analysis must also account for each component: every additional component creates another timeout, schema mismatch, stale index, or verifier false negative to monitor. Conventional microservices already face runtime contract failures across process boundaries. Composed ML systems add probabilistic interfaces in which a large language model planner may hallucinate a tool name or produce JSON that deviates from the declared schema, requiring defensive parsing, retry logic, and output validation at each interface boundary. Capability can increase by adding system structure, but that structure must obey the same latency, reliability, and observability requirements as any production ML system.

System composition aligns naturally with the systems engineering principles studied throughout this book. Modular components can be independently compressed and accelerated using the techniques from chapter 10 and chapter 11. Each component has its own silicon contract (principle 4) and arithmetic intensity profile, allowing hardware-specific optimization. The interfaces between components create natural monitoring points for detecting drift, skew, and degradation. The engineering challenges ahead require mastery across the full stack we have explored: reliable orchestration of multiple models, efficient routing of requests across specialized components, and maintaining consistency across shared state all demand integration from data engineering through model optimization to operational infrastructure.

Systems Perspective 16.1: A new golden age

Hennessy and Patterson (2019) declared a “New Golden Age for Computer Architecture,” driven by the end of Dennard scaling, the slowdown of Moore’s Law, and new opportunities from domain-specific architectures, open instruction sets, and agile chip development. Large-scale AI workloads are one domain where those pressures are visible. Building more capable AI services will not be a matter of writing a better loss function alone; it will be a systems engineering challenge involving energy efficiency, high-bandwidth interconnects, memory

hierarchy design, and software stacks that keep heterogeneous hardware useful rather than idle. The thirteen principles provide a quantitative vocabulary for navigating that regime while preserving the distinction between physical bounds and deployment-specific assumptions.

That era demands concrete engineering advances. Achieving exascale sustained throughput ($\geq 10^{18}$ FLOP/s) and beyond requires new approaches to power delivery, cooling, interconnects, and software coordination, not merely faster chips. The analytical tools developed in this book equip engineers to navigate that regime. Their application evolves as deployment scale and workload mix change.

16.5 Journey Forward

Every frontier just explored rests on a common foundation: the engineering skills this book has developed. Managing stochastic data through versioning and statistical validation, while enforcing execution constraints through physical bounds, runtime checks, and product SLOs, requires bridging the gap between Software 1.0's explicit logic and Software 2.0's learned behavior. Making those assumptions measurable and their failure modes observable is the engineering rigor required to make probabilistic systems dependable.

Intelligence is a systems property. It emerges from integrating data, models, hardware, software, monitoring, and governance rather than from any single breakthrough. The systems lesson is therefore not a recipe for one model family or infrastructure stack. It is the discipline of making every dependency visible enough to measure, every trade-off explicit enough to evaluate, and every deployment responsible enough to operate in the world.

The engineering responsibility


The systems integration perspective explains why ethical considerations cannot be separated from technical ones. Compute requirements affect who can access a system: a model requiring several high-end data center accelerators for inference excludes organizations that cannot afford that infrastructure. Training data can encode biases that affect the system's behavior. Workload energy accounting contributes to data-center carbon footprints that affect the planet. Efficiency choices, data choices, and deployment choices therefore distribute costs and benefits beyond the engineering team. Technical decisions are ethical decisions, viewed through a wider lens.

The question confronting us as engineers is not only what capabilities we can build, but whether we can build those systems well. They must be efficient enough to widen access, secure enough to resist exploitation, sustainable enough to limit environmental harm, and responsible enough to serve people equitably. Systems such as planetary-scale climate monitors and personalized medical assistants require the engineering expertise this book has developed, guided by the responsibility that chapter 15 established as a first-class design constraint.

The principles established here provide a coherent lens for individual ML systems. Larger systems do not invalidate that lens; they expose the same constraints at a different boundary.

A horizon note: From node to fleet

Some workloads eventually exceed a single system. The same bottleneck reasoning still matters, but the resource boundary moves outward. Local memory constraints are joined by network topology constraints, local failure handling becomes fleet reliability, and training throughput becomes a coordination problem. Under an independent, identical constant-hazard model in which the first GPU failure counts as a pool event, the book's reference component mean time to failure (MTTF) of 5.7 years becomes a cluster mean time between failures (MTBF) of about 48.8 hours in a 1,024-GPU pool, before accounting for correlated failures. A run lasting weeks or months therefore has a high probability of encountering hardware failure. A fault-tolerance strategy is an operational requirement for completing such a run efficiently; asynchronous checkpointing and localized recovery can reduce overhead but are not individually required for convergence. That changed boundary is the next frontier: scale. The point is not that this book must become a distributed-systems catalog. The point is that the ML systems lens developed here remains useful when scale changes: identify the binding constraint, quantify the cost term, and trace where that cost propagates.



Fleet scale turns rare component failures into routine system events.

Mastery, however, carries a recurring temptation: the belief that understanding a system means understanding it completely. That temptation produces the fallacies and pitfalls that follow, in which confidence outpaces humility.

16.6 Fallacies and Pitfalls

Fallacies and pitfalls in ML systems arise from a common source: treating the system as decomposable into independent parts. Each fallacy assumes that optimizing one dimension, one metric, or one stage suffices; each pitfall shows the consequence when that assumption meets production reality.

Fallacy: *Systems engineering complexity disappears with better tools and abstractions.*

Tools abstract complexity; they do not eliminate physical constraints. A high-level framework that hides memory management still consumes memory. An AutoML system that tunes hyperparameters still faces the Pareto frontier. Simplifying one interface may shift burden to another, although good design can also remove accidental complexity. The engineer who believes tools eliminate fundamental constraints will be surprised when those constraints resurface at scale, often in forms harder to diagnose than the original problem.

Pitfall: *Optimizing one metric without tracing displaced costs.*

When an optimization reduces latency by 50 percent, ask what changed elsewhere. Quantization may add validation burden. Caching may trade memory capacity for serving speed. Some simplifications remove accidental complexity; others displace cost. Engineers who celebrate gains in one metric without tracing those effects can build systems that fail in unexpected ways. Measurement decides which occurred.

Fallacy: *Mastering individual components equals mastering the system.*

Component expertise is necessary but insufficient. An engineer who understands data pipelines, training, serving, and operations as isolated domains will still struggle with systems where a data schema change cascades through training, breaks quantization assumptions, and triggers silent accuracy degradation in production. Integration can create more complexity than the components reveal in isolation because interfaces multiply failure modes. Systems thinking means understanding how components interact, not just how they work individually.

Pitfall: *Scaling data collection without measuring marginal information value.*

The intuition that more data yields better models is seductive because it often holds early in model development. Chapter 9 demonstrated the diminishing returns that can set in once a dataset achieves sufficient coverage: beyond that threshold, doubling dataset size may yield marginal accuracy gains while increasing storage, preprocessing, and labeling costs. The data-gravity principle recommends measuring those downstream costs, including whether moving or repeatedly scanning a larger dataset is more expensive than moving compute toward it. The engineer who scales data without measuring the incremental return per sample optimizes the wrong variable.

Fallacy: *A single accuracy metric captures model quality.*

A model evaluated solely on accuracy inhabits a one-dimensional world. Pareto analysis includes latency, throughput, memory, energy, fairness, and cost after objective directions are normalized. Under a 100 ms tail-latency SLO, a 95 percent-accurate model at 500 ms is infeasible while a 93 percent-accurate model at 50 ms remains a candidate; without such a requirement, neither point is automatically better. Chapter 15 showed that aggregate accuracy can also conceal large error-rate disparities across demographic groups, so even the accuracy dimension requires disaggregated measurement. Evaluation must span the relevant Pareto surface, not a single axis.

Pitfall: *Treating every drift alarm as an automated rollback trigger.*

Drift detection should trigger a diagnosed response, not an unconditional rollback. Rollback is appropriate when a version or release regression is established. External distribution drift is not repaired by restoring the same stale model; safer automatic actions may include alerting, fallback, traffic reduction, or holding predictions for review, followed by data collection and retraining when outcome evidence supports it. Without cause-specific response mechanisms, a quality regression

can continue until a human notices. Chapter 14 therefore couples monitoring to action, but the action follows the diagnosed cause rather than the drift signal alone.

Fallacy: *A single optimized pipeline stage makes the system fast.*

Amdahl's Law (principle 8) applies directly to end-to-end ML pipelines. Optimizing accelerator inference latency by $10\times$ yields only $1.1\times$ system speedup if CPU-bound preprocessing accounts for 90 percent of end-to-end latency—serial fractions that can arise from host-side image augmentation, synchronous feature-store lookups, or tokenization that remains host bound in some implementations. The iron law of ML systems (principle 3) decomposes execution time into data movement, computation, and latency terms precisely so that engineers can identify the dominant term before investing optimization effort. Chapter 12 formalized this diagnostic process through profiling methodologies that measure where time actually goes. Engineers who optimize without profiling are guessing, and Amdahl's Law is unforgiving of guesses that target the wrong term.

Pitfall: *Profiling only the stage that looks easiest to optimize.*

Teams often profile the model kernel because it is visible, instrumented, and owned by the ML team, while the surrounding data path is split across storage, preprocessing, networking, and application code. That local view can make a $10\times$ kernel improvement look urgent even when it changes little about the user-visible path. End-to-end profiling keeps the optimization target honest: the stage to improve is the one that limits the system, not the one with the cleanest benchmark harness.

All eight fallacies and pitfalls share a common root: the temptation to reduce a system to its parts, whether by optimizing a single metric, a single stage, or a single moment in time. The final summary resists that reduction by returning to the integrated perspective: reasoning across boundaries is the core discipline of ML systems engineering.

16.7 Summary

ML systems engineering differs from isolated component optimization by reasoning across boundaries. The thirteen principles, the conservation-of-complexity heuristic, and the lighthouse journey framework provide analytical tools for reasoning about systems as wholes. Their stated assumptions distinguish exact bounds from fitted models, SLO policies, and design heuristics, allowing the tools to remain useful as frameworks, hardware generations, and model families change.



Key Takeaways: Reasoning across boundaries

- **Assumptions matter across implementations:** The thirteen principles turn framework-specific craft into measurable reasoning by combining physical bounds, decompositions, fitted models, requirements, and heuristics. Apply each only within its stated scope.
- **Trace displaced costs without assuming conservation:** Compression, batching, monitoring, and governance can relocate burdens across data, algorithm, and machine, while good design can remove accidental complexity outright. Measurement distinguishes the two cases.
- **Boundaries reveal the bottleneck:** In the idealized batch-one, two-H100 model, the memory-transfer lower bound for a Llama 2 70B token is about $295.2\times$ the peak-compute lower bound, and the illustrative p99 latency is $40\times$ the mean. Systems thinking means measuring where physics, traffic, and users bind.
- **Scale changes the binding term:** The next frontier is scale, where a thousand-GPU pool turns multi-year component MTTF into days-scale cluster MTBF. The physics stays, but the constraint moves to fleets.

Hennessy and Patterson's *Computer Architecture: A Quantitative Approach* helped establish a shared analytical language for comparing CPI, clock rates, and instruction counts (Hennessy and Patterson 2011; Hennessy and Patterson 2017). This collection aspires to a similar role for ML systems engineering without claiming that every diagnostic is an invariant. It is a beginning, not an endpoint. Future work will refine the models, assumptions, and scope.



What will endure is the intellectual posture these principles embody: reasoning from evidence and physical bounds rather than reacting to symptoms, quantifying trade-offs rather than following trends, and treating design as constrained optimization. This is the engineering corollary of the bitter lesson the introduction drew from seven decades of AI research: because general methods that scale with computation have repeatedly outrun hand-crafted expertise, the durable advantage belongs to systems engineering that can absorb that computation, not to any single clever architecture (Sutton 2019). Specific frameworks will rise and fall, hardware generations will turn over, and model architectures will be superseded. Disciplined reasoning about data, computation, and physical constraints will not.

At the next frontier of scale, some models no longer fit on one machine, failures become highly probable across fleets, and network links can become binding alongside local memory buses. The physics does not change; the scale at which it binds does.

The world is rushing to build AI systems. Our task is to engineer them.

Prof. Vijay Janapa Reddi, Harvard University

BACK MATTER



The D·A·M Taxonomy

A system that runs slowly rarely announces why. An idle accelerator can stem from a starved input pipeline (Data), an algorithm with too little arithmetic intensity (Algorithm), or a saturated memory bus (Machine). This appendix formalizes the book’s single-node diagnostic taxonomy, mapping profiler symptoms and measurements directly to the axis that actually binds execution before optimization begins. Readers should arrive with the three axes from the opening chapters and the iron law that quantifies them.

How to Use This Appendix

Use the D·A·M taxonomy when a system misses a performance, cost, or quality target and the cause is not yet clear. Work through the framework in four passes:

1. **State the symptom and operating budget:** Record the missed end-to-end target, workload, and deployment context.
2. **Localize the candidate axis or boundary:** Use table A.1 for a first pass and table A.2 when the evidence crosses axes.
3. **Test the hypothesis:** Map the evidence to the iron law and arithmetic intensity, then gather confirming measurements with table A.6 and table A.7.
4. **Intervene and remeasure:** Choose a lever from table A.4 and repeat the end-to-end measurement. If the target does not improve, revise the hypothesis.

The symptom determines the most useful entry point. Low accelerator utilization or slow batch delivery should begin with section A.5 and section A.8. A missed service-level objective (SLO) should move from phase timing and queue depth in table A.5 to the corresponding iron-law term. A quality plateau should be checked against the Data/Algorithm boundary and section A.11 before allocating more compute. In every case, the axis label is a hypothesis, not a verdict: the axes overlap, and a useful diagnosis may name both a dominant axis and a boundary effect.

A.1 Diagnostic Summary

The taxonomy maps directly to the iron law of ML systems established in section 1.7. Table A.1 summarizes the role, primary physical constraint, and core optimization pathway for each axis.

Table A.1: D·A·M Axis Reference: Each axis maps to a constraint class and optimization strategy. Form a hypothesis, confirm it with measurements, then follow the chapter pointer.

Axis	Role	Physical Constraint	High-Leverage Optimization
Data (D)	Information (The Fuel)	Volume, quality, arrival rate	Data Selection (chapter 9)
Algorithm (A)	Logic (The Blueprint)	Operations and dependencies	Model Compression (chapter 10)
Machine (M)	Physics (The Engine)	Compute, memory, I/O	Hardware Acceleration (chapter 11)

This first-pass separation is useful, but production systems rarely suffer from a single axis-centered bottleneck. More often, the problem sits at the *boundary* between two axes—a data format choice that determines whether the GPU can be saturated, or a pruning strategy that changes the memory access pattern. Handling these cases requires mapping the intersections.

A.2 Intersection Landscape

The opening chapter uses figure 1.4 to establish the conceptual landscape. Here, figure A.1 turns that map into a diagnostic sequence. First identify the binding axis, then ask how an intervention will alter the other two.




Figure A.1: The D-A-M Axis Triad: A compact diagnostic view of Data, Algorithm, and Machine as three interdependent axes. The bidirectional links emphasize that an intervention on one axis changes the constraints on the other two.

With the first-pass axis identified, table A.2 expands the pairwise and three-way intersections into concrete engineering techniques and their primary book chapters.

Table A.2: D-A-M Intersection Reference: Each zone maps specific techniques to the axes they span and the chapters that cover them. The pairwise intersections require reasoning about two domains simultaneously; the center requires all three.

Zone	Name	Key Techniques	Book Coverage
Data	Information	Storage formats, data quality, distributions	chapter 4
Algorithm	Logic	Loss functions, architectures, gradients	chapter 5, chapter 6
Data $\cap$ Algorithm	What to Learn From	Data selection, curriculum learning, compute-optimal scaling	chapter 9, chapter 8
Data $\cap$ Machine	How to Move Information	I/O bandwidth, prefetching, data formats	chapter 4, chapter 11
Algorithm $\cap$ Machine	How to Execute Efficiently	Quantization, pruning, kernel fusion, mixed precision	chapter 7, chapter 10
Machine	Physics	Silicon, memory hierarchy, peak FLOP/s	chapter 11
Data $\cap$ Algorithm $\cap$ Machine	ML Systems Engineering	Iron law, Roofline, training loops, serving	chapter 8, chapter 13, chapter 12

In the D-A-M acronym, Data, Algorithm, and Machine are taxonomy axes. They are not mathematical variables; formal quantities still follow the notation chapter, where D denotes dataset size or training tokens.

The axis-centered zones contain concepts that belong primarily to one axis: storage formats and distributions emphasize Data, loss functions and gradients emphasize Algorithm, and silicon physics and peak FLOP/s emphasize Machine. Single-domain expertise can introduce them, but deployed behavior may still depend on the other axes.

The pairwise intersections are where systems thinking begins. Data $\cap$ Algorithm (*What to Learn From*) encompasses data selection, curriculum learning, active learning, and scaling laws like Chinchilla ($D \approx 20P$ for the dense autoregressive language-model regime studied) (Hoffmann et al. 2022)—all requiring joint reasoning about information content and algorithmic capacity. Adding data without considering whether the model can learn from it wastes compute; choosing architectures without considering data availability wastes engineering time. Data $\cap$ Machine (*How to Move Information*) covers I/O bandwidth, prefetching strategies, data formats, and the energy-movement invariant. This intersection is where data gravity manifests: the physical cost of moving bytes through the memory hierarchy determines whether the machine can be fed fast enough. Algorithm $\cap$ Machine (*How to Execute Efficiently*) spans quantization, pruning, kernel fusion, mixed precision, and computational graph optimization. A pruning strategy that reduces FLOPs but destroys memory access patterns can *slow down* execution on real hardware.

The center—Data $\cap$ Algorithm $\cap$ Machine—is where all three axes converge. The iron law, the Roofline Model, end-to-end training loops, serving pipelines, and holistic benchmarking all require simultaneous reasoning about data flow, algorithmic complexity, and hardware utilization. This center is not a single technique; it is the discipline itself.

Understanding the landscape reveals *where* a technique lives. The next step is quantifying *which axis dominates* for a given workload—and that requires the iron law.

A.3 Iron Law Mapping

For serialized phases, the iron law maps work across the D·A·M axes to execution time. Here, T and L_{lat} are times in seconds; D_{vol} is bytes moved and BW is bytes per second; O is work in FLOPs, R_{peak} is FLOP/s, and η_{hw} is dimensionless hardware efficiency. These quantities map to the axes as follows:

$$T = \underbrace{\frac{D_{\text{vol}}}{\text{BW}}}_{\text{Data/Machine}} + \underbrace{\frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}}_{\text{Algorithm/Machine}} + \underbrace{L_{\text{lat}}}_{\text{Cross-axis overhead}}$$

Algorithm and Machine share the compute term, separated by which variable the engineer controls. Reducing the total operations (O) is an Algorithm lever, while improving the hardware’s peak throughput (R_{peak}) or utilization (η_{hw}) is a Machine lever.

This equation transforms performance debugging from a qualitative guessing game into a quantitative engineering problem. The dominant measured cost appears in one of these terms. A slow system may move too much data (D_{vol}), lack bandwidth (BW), execute too many operations (O), fail to use the hardware’s peak capability (η_{hw}), or pay cross-axis overhead (L_{lat}). The levers below map specific optimizations to the variables they improve.

A.3.1 Component levers

- **Data Lever:** Reducing moved bytes (D_{vol}) through deduplication, selection, or lower-precision representation, or increasing I/O bandwidth (BW).
- **Algorithm Lever:** Reducing operations (O) through pruning or architectural refinement. Quantization narrows representation and may raise hardware throughput; it does not generally reduce operation count.
- **Machine Lever:** Increasing the denominator of the compute term by improving peak throughput (R_{peak}) or increasing the utilization factor (η_{hw}) via kernel fusion.

A.3.2 D·A·M coordination: From sum to max

The additive iron law represents sequential execution. Let $T_{\text{sequential}}$ denote time without overlap and $T_{\text{pipelined}}$ denote measured time with overlap. Overlap can move the sum toward a max, but dependencies, contention, and incomplete overlap keep the measured time above this ideal lower bound:

$$T_{\text{sequential}} = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}} \xrightarrow{\text{overlap}} T_{\text{pipelined}} \geq \max \left(\frac{D_{\text{vol}}}{\text{BW}}, \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} \right) + L_{\text{lat}}$$

The systems engineer’s job is to make these components run in parallel, not in series. Table A.3 summarizes key D·A·M coordination techniques that overlap work across axes:

Table A.3: D·A·M Overlap Techniques: Each technique overlaps work across axes, moving sequentially added Data/Machine and Algorithm/Machine time toward their ideal lower bound: the larger of the two terms. Dependencies and contention limit realized overlap.

Technique	D·A·M Axes Overlapped	Implementation
Prefetching	D overlaps M	DataLoader with <code>prefetch_factor</code> , <code>pin_memory=True</code>
CUDA Streams	D overlaps M	Separate streams for H2D transfer and compute

Technique	D·A·M Axes Overlapped	Implementation
Async Gradient Sync	M (communication) overlaps A	Overlap bucketed AllReduce with remaining backward computation
Double Buffering	D overlaps M	Fill buffer N+1 while computing on buffer N

Overlaps provide the greatest relative speedup when the D·A·M time terms are reasonably balanced. If one term dominates (for example, severely memory bound), overlapping the smaller term with the larger yields negligible gain—the max is still dominated by the same bottleneck. This balance occurs near $D_{\text{vol}}/\text{BW} \approx O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$. The latency term requires separate treatment.

🔍 Systems Perspective A.1: The overhead that cannot hide

The unhidden latency term L_{lat} (kernel launch, synchronization barriers, Python dispatch) contains serialization points that cannot be fully overlapped. Kernel fusion can reduce this term by combining operations, but it does not eliminate all launch, scheduling, or synchronization overhead.

The iron law tells the engineer *how much* time each term consumes. One critical question remains: does data movement or arithmetic throughput set the active ceiling? The answer lies in a single ratio.

A.4 Arithmetic Intensity Boundary

The boundary between the *Data/Machine* memory term and the *Algorithm/Machine* compute term is not arbitrary; it is defined mathematically by arithmetic intensity¹ (I) of the workload.

The Roofline Model provides an upper performance bound when traffic and achieved rates are measured consistently. In the middle of a production incident, faster screening heuristics can point to the measurements to collect next.

A.5 Rules of Thumb

In the heat of a production outage, there is rarely time to solve the full iron law equation. Veteran systems engineers instead rely on quantitative heuristics to quickly narrow the search space; the thresholds below are screening signals that should be checked against profiler traces and hardware counters.

- Low accelerator utilization (< 80 percent): Screen for data, CPU, or launch starvation. Confirm with input-pipeline wait, host CPU saturation, memory-bandwidth counters, and trace gaps.
- High accelerator utilization (> 95 percent): Treat this as machine bound only if compute units are busy and memory bandwidth is not saturated. If memory bandwidth is saturated, the bottleneck is still on the Data/Machine boundary.
- If batch size is one: Treat this as a warning that latency or launch overhead may matter; confirm with traces before labeling the algorithm itself the bottleneck.
- Low arithmetic intensity (below the hardware ridge point): The workload is likely memory bound (Data/Machine boundary). The precise boundary is $R_{\text{peak}}/\text{BW}$ and depends on hardware and precision.
- If the system works in dev but fails in prod: Treat data drift as one hypothesis alongside configuration, load, dependency, and runtime differences.

Common industry labels cross the axes. Memory-bound execution sits on the Data/Machine boundary because both bytes moved and available bandwidth matter. Compute-bound execution couples Algorithm work with Machine throughput. Latency-bound execution may reflect algorithmic serial depth, runtime dispatch, synchronization, or queueing.

¹ | **Arithmetic intensity:** This book uses the term for floating-point operations per byte transferred. Williams et al. (2009) calls the measured DRAM quantity *operational intensity* in the Roofline Model. Comparing the workload ratio with the hardware ridge point ($R_{\text{peak}}/\text{BW}$) gives an upper-bound diagnosis of memory- versus compute-bound execution.

A.5.1 Bottleneck diagnostic

Once profiling identifies the dominant term in the iron law, optimizations must target that specific axis. Table A.4 contrasts the interventions that yield end-to-end speedups against nonbinding optimizations that waste engineering effort:

Table A.4: What Works vs. What Is Wasted: Optimizing a nonbinding term yields little or no end-to-end improvement. A memory-bound large language model will not benefit from additional peak FLOP/s alone, though an accelerator with more memory bandwidth may help.

If the workload is...	Dominant Term	Optimization That Works	Optimization That is Wasted
Memory-Bound	D_{vol}/BW	Quantization; structured pruning; batching for weight reuse; fusion that removes intermediate traffic	Faster accelerator (more FLOP/s will not help)
Compute-Bound	$O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$	Better kernels, Tensor Cores, faster accelerator, lower precision	More memory bandwidth (not the binding term)
Latency-Bound	L_{lat}	Reduce serial depth; fuse or dispatch asynchronously; batch within budget	More compute or bandwidth unless the trace shows it binding

Knowing what works also means recognizing what does not. In practice, teams under deadline pressure repeatedly fall into the same traps—optimizing the wrong axis with confidence. These failure modes are common enough to deserve their own names.

A.6 Anti-Patterns

Diagnosing systems is often a process of elimination. Before committing to complex kernel optimizations, watch for these common traps that waste engineering cycles.

- The hardware crutch: Buying faster accelerators (*Machine*) to fix a slow Python data loader (*Data*). The new hardware will just idle faster.
- The model twiddle: Changing neural architectures (*Algorithm*) when the bottleneck is actually network bandwidth or disk I/O.
- The premature optimizer: Writing custom CUDA kernels (*Machine*) before verifying if the Algorithm is simply doing too many unnecessary operations.

Each anti-pattern follows the same root cause: acting before diagnosing. The following case studies show what proper diagnosis looks like—starting from a confusing symptom and systematically narrowing to the dominant D·A·M axis.

A.7 D·A·M Case Studies

Theoretical constraints often manifest as confusing symptoms in production. These representative scenarios illustrate how to apply the taxonomy. Each case follows the same three diagnostic moves: symptom, diagnosis, and fix.

A.7.1 Case 1: The starving accelerator (Data)

A team provisions a large A100 GPU instance to speed up training, but training time hardly improves and `nvidia-smi` shows GPU utilization fluctuating between 10 percent and 40 percent. The pattern suggests accelerator starvation but does not identify its cause. If traces show input-wait gaps with storage or CPU decoding saturation, the data path is binding. The useful intervention is then upstream of the model. Optimize the extract, transform, load (ETL) path by moving from raw JPEGs with heavy CPU decoding to sequential formats such as TFRecords or WebDataset, increasing data-loader parallelism, and prefetching batches into accelerator memory.

A.7.2 Case 2: The latency cliff (cross-axis)

A real-time recommendation service fails to meet a 20 ms latency service level agreement (SLA), while accelerator utilization is low and the batch size is one. That combination suggests launch overhead, memory traffic, or serial model depth rather than a saturated chip. If a trace confirms

that the sequential layer path dominates, adding more hardware will not remove it, so pruning or knowledge distillation can reduce model work. If weight traffic dominates instead, quantization can reduce moved bytes by using INT8 where accuracy allows. The measurements select the remedy.

A.7.3 Case 3: The compute wall (Machine)

Accelerator utilization is pinned at 99 percent, memory bandwidth remains unsaturated, and training is stable but takes three weeks. If profiler counters also show high compute-pipeline occupancy and model FLOPs utilization (MFU), the workload is compute bound. The data path is keeping the accelerator fed, so the next step must change the compute term in the iron law. The team can scale up from an A100 to an H100-class accelerator, scale out through data parallelism, or lower precision from FP32/TF32 to BF16 where numerically safe. On NVIDIA Tensor Core paths, BF16 peak throughput is typically about $2\times$ TF32 peak, with realized speedup depending on kernels and bottlenecks (Choquette 2023). Taken together, the three cases show why the same utilization number can imply different next steps depending on loss behavior, batch size, and memory pressure.



Checkpoint A.1: D·A·M diagnosis check

1. A training job shows 95 percent accelerator utilization but loss has plateaued for two epochs. Which D·A·M axes could be responsible, and why is utilization alone insufficient to distinguish them?
2. Your colleague suggests adding more data loader workers to a job where `nvidia-smi` shows 98 percent GPU utilization. What trace evidence is needed before deciding whether this will help?
3. An inference server meets its latency SLO at batch size 1 but fails at batch size 16. Which iron-law terms and queueing effects may have changed, and which measurements identify the binding regime?

These three cases illustrate bottlenecks that can cross axes. Production incidents are rarely so tidy—symptoms often overlap, and the dominant axis can shift during debugging. The next section provides a systematic troubleshooting matrix for the messier scenarios encountered in practice.

A.8 Production Troubleshooting

Identifying the root cause of performance bottlenecks requires systematic elimination. Table A.5 provides a diagnostic matrix for common failure modes observed in production deployments.

The diagnostic matrix indicates *what* to suspect. The next question is *how* to confirm that suspicion with evidence—which requires the right profiling tools.

Table A.5: D·A·M Diagnostic Matrix: First hypotheses and measurements for common failures. Confirm the binding component before intervening.

Symptom	Hypothesis	Confirm With	Act After Confirmation
Low Accelerator Utilization	Data/host starvation	Trace gaps; loader, CPU, storage, launch counters	Prefetch, parallelize, or cut launches.
High Latency (P99)	Algorithm/run-time/queue	Phase p99 and queue depth	Optimize the dominant phase.
High Training Cost	Low useful-work efficiency	MFU, communication, idle time	Reduce the largest loss or resize.
High Inter-GPU Overhead	Communication bound	Device trace; collective gaps	Bucket AllReduce; tune data-parallel overlap.
Out-of-Memory (OOM)	Algorithm/Machine fit	Peak weights, states, activations	Checkpoint, shard, quantize, or shrink batch.

A.9 Tooling Map

Once a hypothesis exists (for example, “the workload appears Machine-bound”), evidence is needed to confirm it. Abstract concepts must be measured with concrete utilities. Table A.6 connects the theoretical components to the specific Linux and Python profiling tools that confirm them.

Table A.6: D·A·M Tooling Map: Profiling utilities for diagnosing bottlenecks along each D·A·M axis. Start with the primary tool for quick triage; use secondary tools for deep-dive analysis when the primary tool’s output is inconclusive.

Axis	Key Metric	Primary Tool	Secondary Tool
Data	Batch Load and Wait Time	Framework profiler trace	iostat, pidstat (I/O/CPU)
Algorithm	FLOPs, Model Depth	PyTorch Profiler	DeepSpeed Flops Profiler
Machine	Accelerator Utilization, SM Occupancy	Nsight Compute	nvidia-smi, Nsight Systems

Profiling tools generate raw numbers—utilization percentages, FLOP counts, and bandwidth measurements. They become actionable only in workload and hardware context. The D·A·M Scorecard provides screening signals rather than universal standards, narrowing which measurements and interventions deserve attention.

A.10 D·A·M Scorecard

To move beyond qualitative guessing, the efficiency indicators in table A.7 help decide what to investigate next. This report-card view aids the first comparison, but no row provides a universal grade. MFU² is especially useful for large-model training.

Table A.7: D·A·M Screening Indicators: Context-dependent signals for investigation, not universal grades.

Axis	Metric	Definition	Investigate	Strong Signal
Data	I/O Overhead	$\frac{\text{Data Wait Time}}{\text{Total Step Time}}$	> 10% (screen)	< 1% (strong)
Algorithm	Avoidable Work	Profiler or ablation finds redundant operations	Redundant work found	End-to-end work falls
Machine	MFU	$\frac{\text{Achieved model FLOP/s}}{\text{Peak FLOP/s}}$	< 30% (large models)	> 50% (same regime)

The Scorecard and the Roofline Model both answer efficiency questions, but at different scales. The Scorecard screens the current system using context-dependent indicators. Scaling laws and the information-roofline metaphor ask what may happen as the system scales *beyond* its current size.

A.11 Scaling Laws vs. Roofline

Systems engineering requires distinguishing between *growth trajectories* and *fundamental limits*.

A.11.1 Scaling laws (the journey)

Scaling laws³ are empirical power laws that predict how held-out loss changes as resources increase within the fitted regime. Two landmark results are Kaplan scaling (Kaplan et al. 2020), which studied loss against parameter count (P), data (D), and total operations (O), and Chinchilla scaling (Hoffmann et al. 2022), which estimated a compute-optimal balance for the dense autoregressive language-model regime studied, often summarized as roughly 20 training tokens per parameter ($D \approx 20P$).

As economic guides, these fits estimate how held-out loss changes with compute inside the observed regime. They do not guarantee a task error-rate reduction or extrapolate unchanged across architectures and data distributions.

² | **MFU (model FLOPs utilization):** The ratio of achieved model FLOP/s to the hardware’s theoretical peak FLOP/s, introduced in the PaLM paper (Chowdhery et al. 2022). Unlike raw accelerator utilization (which counts any work the accelerator performs), MFU measures model computation rate relative to peak hardware throughput, so non-model work and system overhead are not counted as achieved model FLOPs. Section 8.5 develops MFU as a training diagnostic.

³ | **Scaling laws:** Empirical relationships, typically power laws with a fitted negative exponent, that predict model loss as a function of dataset size, parameter count, or compute budget. Kaplan et al. (2020) studied these relationships for neural language models at OpenAI; Hoffmann et al. (2022) later estimated the compute-optimal training trade-off for its studied dense language-model regime. Section 9.11.2 develops the compute-optimal frontier and its data-starvation diagnostic in detail.

A.11.2 Information roofline (the destination)

The *information roofline* is used here as a diagnostic metaphor rather than a standard quantitative law. Its destination image points to limits on what can be learned from the available data. For classification under a fixed data distribution and loss, the Bayes error rate⁴ is an irreducible error *floor*, not a ceiling. Noise, label ambiguity, coverage gaps, and distribution shift can create data-quality limits, but they do not define one universal slope or breakpoint.

The diagnostic lesson is narrower. Scaling laws describe a fitted improvement trajectory, while the information-roofline metaphor prompts a search for data-quality limits. If a loss curve flattens relative to the fitted trend, investigate optimization, capacity, evaluation noise, and data quality; flattening alone does not identify which cause dominates. The Bayes floor describes the best achievable classifier for the specified distribution. Adding accelerators or parameters is futile only after measurements establish that data quality is binding. Whether debugging a training step, evaluating hardware utilization, or planning a scaling campaign, measure which axis is binding before choosing a lever.

A.12 Summary

These tools form a single diagnostic workflow. D-A-M localizes candidate causes across Data, Algorithm, and Machine. The iron law and arithmetic intensity translate those causes into time, operations, and bytes, while the case studies, anti-patterns, tooling map, and scorecard show how to test them in production. Scaling laws and the information roofline then separate observed improvement trajectories from data-quality limits. Use the framework as a loop rather than as a label.



Key Takeaways: Where to look first

- Begin with cross-axis hypotheses; a symptom rarely identifies its cause.
- Translate each hypothesis into the relevant iron-law term and a measurement that could falsify it.
- Use arithmetic intensity and the scorecard to identify the binding constraint, not to assign universal grades.
- Apply the corresponding lever and remeasure end to end; a local gain that leaves the binding term unchanged is not a system improvement.

⁴ | **Bayes error rate:** The lowest achievable error rate under 0–1 loss for any classifier on a given data distribution, determined by class priors and the overlap between class-conditional distributions (Goodfellow et al. 2016). No amount of data, parameters, or compute can reduce error below this theoretical floor.

Data Foundations

DATA problems rarely present as data problems; they surface as a training run that will not speed up, an accelerator sitting idle, or accuracy eroding while dashboards stay green. To keep chapter arguments focused on system design, this appendix gathers the back-of-the-envelope calculations and statistical tools that trace such symptoms back to measurable data bottlenecks: transfer times, storage-layout costs, drift diagnostics, and numerical-stability tools. The tools here lean on basic probability and on the data term of the iron law.

How to Use This Appendix

This appendix is designed as a *reference*. Reach for it when debugging “slow training,” “low accelerator utilization,” or “production accuracy drift” calls for the quickest path from symptom to measurement.

Conventions used here follow the book-wide notation (for example, B is reserved for batch size and BW denotes bandwidth).

- **When data will not move:** Start with table B.1 and the transfer-time equation in section B.1.1.
- **When the accelerator is starving:** Use table B.2 and the layout discussion in section B.1.3.
- **When pipelines explode in cost:** Use the primitives in section B.1.4, especially join-induced shuffles.
- **When “average looks fine” but users complain:** Use section B.2.1 and session-level tail probability.
- **When accuracy drifts silently:** Use section B.2.2 to compare full distributions.

The data landscape can feel like a zoo of formats, encodings, and edge cases, but this appendix focuses on the structural constants that constrain every ML system.

With those reference points in place, the appendix begins with the physical constraints of data engineering and then moves to the statistical monitoring that keeps ML systems healthy. From storage formats that determine I/O throughput to drift metrics that detect silent failures, these foundations connect directly to the data pipelines in section 4.5 and the operational monitoring in section 14.5.2.2.

B.1 Data Engineering Foundations

Understanding hardware constraints is only half the battle; we must also shape our data to fit them. Data engineering applies the principles of the memory hierarchy to storage formats and pipeline design, ensuring that the accelerator never starves. This process begins with recognizing that data is physical—it has volume, it takes time to move, and it requires energy to parse.

B.1.1 Napkin math: The physics of data gravity

Data gravity¹ is a metaphor grounded in transfer-time calculations. Unlike compute, which gets faster every year, the speed of light is fixed and network bandwidth is a finite resource. When datasets grow large enough, their practical inertia grows—moving them can cost more time and energy than moving computation to where the data already lives.

The ideal line-rate transfer time is $T = D_{\text{vol}}/\text{BW}$ after converting data volume and bandwidth to compatible units (the calculations below convert decimal bytes and bits per second explicitly). For

¹ | **Data gravity:** Coined by Dave McCrory in 2010 to describe how large datasets attract services and applications toward them, much as massive bodies attract smaller ones in physics (McCrory 2010). The analogy is apt: the “escape velocity” required to move a petabyte-scale dataset is often measured in weeks.

the large volumes here, latency is negligible; the full equation appears in section D.3.4. Table B.1 gives lower bounds; protocol overhead and endpoint limits increase them.

Table B.1: The Cost of Inertia: These transfer times explain why we often “ship compute to data” rather than “ship data to compute.” Network columns are lower bounds at sustained line rate. The truck entry is an illustrative elapsed-time model for 1 PB: 8 hours to load, 32 hours in transit, and 8 hours to unload, excluding scheduling and provisioning overhead.

Data Volume	1 Gbps (Standard WAN)	10 Gbps (High-End WAN)	100 Gbps (Direct Connect)	Truck (Illustrative)
1 TB	2.2 hours	13.3 minutes	80 seconds	N/A
100 TB	9 days	22.2 hours	2.2 hours	N/A
1 PB	3 months	9 days	22.2 hours	2 days

B.1.2 The cost of serialization

Even after data arrives at the machine, the serialization tax remains a final hurdle.² As table B.2 shows, many engineers meticulously optimize their accelerator kernels while ignoring the CPU overhead of decoding data. Parsing text-based formats like JavaScript Object Notation (JSON) or CSV is extremely CPU-intensive, often leaving the accelerator idling while the CPU struggles to convert strings into floating-point numbers.

Table B.2: Serialization Overhead: These illustrative values show why binary columnar formats can outperform row-based text, but they are not portable benchmarks. Arrow is an in-memory/IPC format; Parquet is a compressed storage format.

Format	Decoding Speed (MB/s)	Relative CPU Decode Cost	Suitability
CSV/JSON	~100 MB/s	High	Inspection/interchange
Protobuf	~300 MB/s	Medium	RPC/Messages
Parquet	~1,000 MB/s	Low	Columnar storage

B.1.3 Row vs. columnar formats

The choice of file format determines the “physics” of how the data is read. Row-oriented formats such as CSV and JSON store data record-by-record, so reading the age field still requires parsing each row. This suits appended logs but not analytics on selected features. Parquet stores compressed columnar chunks on disk, while Arrow provides a columnar in-memory and IPC format. Parquet readers can use *projection pushdown* to fetch selected column chunks, and both formats support *vectorized processing*. Compare the two arrangements in figure B.1 to see why columnar access avoids scanning unnecessary bytes.




Figure B.1: Storage Layouts: Row-oriented formats group values by record, so the illustrated scan reads every field. Column-oriented formats group values by feature, so the illustrated scan reads only Age.

The layouts in figure B.1 differ in how many irrelevant bytes enter the training path.

Systems Perspective B.1: The accelerator starvation problem

The choice of file format can determine whether a system is I/O bound or compute bound. As figure B.1 shows, storage layout compounds the serialization tax because row-oriented formats force full-row scans while columnar formats enable projection pushdown, reading only the bytes the model needs. The result is that the “Data Movement” term in the iron law can silently become the bottleneck that leaves expensive accelerators idling.

B.1.4 The algebra of data

Feature engineering turns raw records into the columns a model can learn from, and that transformation is usually a dataflow problem before it is a modeling problem. A pipeline first narrows the population, then chooses the features, then attaches context from other tables. Those three moves correspond to three Structured Query Language (SQL) primitives, and their computational cost determines whether the feature job stays local, scans unnecessary bytes, or turns into a network shuffle.

1. **Selection (σ):** Filtering rows (for example, `WHERE age > 30`).
 - **Cost:** $\mathcal{O}(\log N + K)$ if a B-tree index lookup returns K rows; $\mathcal{O}(N)$ for a full scan.
2. **Projection (π):** Selecting columns (for example, `SELECT age`).
 - **Cost:** Selected chunks plus decoding and metadata in columnar formats. This row-format example reads 1 KB for a 4-byte integer, wasting 99.6 percent of the read.
3. **Join ($\bowtie$):** Combining tables. The most expensive operation.
 - **Shuffle Join:** Both tables are partitioned by key and exchanged over the network.
 - **Cost:** Reshuffling both 1 TB tables moves about ~2 TB; partitioning and locality can reduce this.
 - **Broadcast Join:** One small table is sent to all workers.
 - **Cost:** $S(W - 1)$ bytes; the S -byte table must fit on each of W workers.

Understanding data formats, serialization costs, and algebraic primitives tells us how to *move* data efficiently. Even a perfectly engineered pipeline, however, can silently fail if the data it carries changes character over time. Detecting that change—and quantifying how much it matters—requires a different set of tools: probability and statistics.

B.2 Probability and Statistics

Once data is flowing through our pipelines, we need mathematical tools to ensure its quality and consistency. Probability and statistics provide the language for monitoring system health, detecting the silent failures of data drift, and managing uncertainty.

B.2.1 Distributions and the long tail

In systems, the mean is often misleading. Latency distributions are often long-tailed, and user-visible services are commonly governed by high-percentile behavior rather than mean response time (Dean and Barroso 2013). A “P99” (99th percentile) latency of 500 ms means at most 1 percent of requests take longer than 500 ms; with many requests per session, the fraction of users who see at least one slow request can be much higher. At scale (1M users), even a 1 percent affected-user rate would affect 10,000 users.

Napkin Math B.1: The median experience

For $N_{\text{session}} = 100$ independent requests, each with a 0.01 chance of being slow, the probability of at least one P99-tail event is $\Pr(\text{Slow}) = 1 - (0.99)^{100} \approx 1 - 0.366 = 63.4\%$.

Systems insight: Under this model, a 100-request session is more likely than not to include an event beyond P99. Correlated slowdowns can change this probability (Dean and Barroso 2013).

B.2.2 Measuring drift (divergence)

Because distributions have long tails where the most dangerous failures hide, simple metrics like “mean shift” are insufficient. We need tools that compare the *entire shape* of the distribution. Detecting drift between the serving distribution (P_t) and training distribution (P_0) requires measuring the “distance” between distributions.

For discrete outcomes x , equation B.1 expresses KL divergence³ as the expected extra coding cost of approximating P_t with P_0 (Kullback and Leibler 1951; Cover and Thomas 2006). The displayed form is finite only when $P_0(x) > 0$ wherever $P_t(x) > 0$; production implementations therefore need explicit smoothing or zero handling.

$$\mathcal{D}_{\text{KL}}(P_t \| P_0) = \sum_x P_t(x) \log \frac{P_t(x)}{P_0(x)} \quad (\text{B.1})$$



Napkin Math B.2: Worked example: KL divergence for drift detection

Scenario: A sentiment classifier was trained on data where 60 percent of reviews were positive, 30 percent negative, and 10 percent neutral. After deployment, the serving distribution shifts to 45 percent positive, 40 percent negative, and 15 percent neutral.

Math:

Training distribution P_0 : [0.60, 0.30, 0.10]. Serving distribution P_t : [0.45, 0.40, 0.15].

$$\begin{aligned} \mathcal{D}_{\text{KL}}(P_t \| P_0) &= 0.45 \log \frac{0.45}{0.60} + 0.40 \log \frac{0.40}{0.30} + 0.15 \log \frac{0.15}{0.10} \\ &= 0.45 \times (-0.2877) + 0.40 \times (0.2877) + 0.15 \times (0.4055) \\ &= -0.1295 + 0.1151 + 0.0608 = 0.0464 \text{ nats} \end{aligned}$$

$$\begin{aligned} \mathcal{D}_{\text{KL}}(P_0 \| P_t) &= 0.60 \log \frac{0.60}{0.45} + 0.30 \log \frac{0.30}{0.40} + 0.10 \log \frac{0.10}{0.15} \\ &= 0.60 \times 0.2877 + 0.30 \times (-0.2877) + 0.10 \times (-0.4055) \\ &= 0.1726 + (-0.0863) + (-0.0405) = 0.0458 \text{ nats} \end{aligned}$$

Notice the asymmetry: $\mathcal{D}_{\text{KL}}(P_0 \| P_t) \neq \mathcal{D}_{\text{KL}}(P_t \| P_0)$. The population stability index (PSI) symmetrizes this; the calculation below uses natural logarithms:

$$\begin{aligned} \text{PSI} &= \sum (P_{0,i} - P_{t,i}) \log \frac{P_{0,i}}{P_{t,i}} = (0.15)(0.2877) + (-0.10)(-0.2877) + (-0.05)(-0.4055) \\ &= 0.04315 + 0.02877 + 0.02027 = 0.0922 \end{aligned}$$

Systems insight: Since $\text{PSI} = 0.0922 < 0.2$, this drift does not yet cross the illustrative manual-review threshold used here. The right response is to keep the feature on the monitoring dashboard and look for corroborating drift signals rather than retraining from this statistic alone.

The threshold remains illustrative: its meaning depends on sample size, binning, smoothing, and zero handling, and retraining requires model-quality evidence.

Both KL divergence and PSI are grounded in a deeper framework—information theory—which provides the units and bounds that make these metrics principled rather than ad hoc.

³ | **KL divergence:** Named after Solomon Kullback and Richard Leibler, who introduced it in 1951. Also called *relative entropy*, it quantifies the expected extra information needed to encode samples from one distribution using a code optimized for another; for drift monitoring in this book, the canonical direction is $\mathcal{D}_{\text{KL}}(P_t \| P_0)$. With natural logarithms, as in the example above, the unit is nats; with base-2 logarithms, the unit is bits. In production ML, KL divergence is the theoretical backbone of many drift-detection metrics.

B.2.3 Information theory for systems

A training run can consume more examples, run longer, and still stop improving if the added data carries too little useful signal. At that point, the bottleneck is not only compute or bandwidth; it is the amount of information the data pipeline delivers to the learner. Section A.11.2 treats that data quality limit as a diagnostic constraint, and information theory provides the units for reasoning about it.

Entropy (H)⁴ is the average uncertainty in a distribution, defined as $H(X) = -\sum p(x) \log p(x)$. The log base sets the unit: natural logs give nats, while $\log_2$ gives bits. A uniform distribution has maximum entropy only over a fixed finite support. This connects directly to KL divergence above: $\mathcal{D}_{\text{KL}}$ measures the excess coding cost of using the wrong distribution.

Information density is used here as a book-level heuristic for useful learning signal per unit of storage, not as a standard information-theory quantity.

Signal-to-noise ratio (SNR) is the ratio of signal power to noise power. In ML, the book uses low SNR qualitatively for data in which task-relevant structure is weak relative to noise; adding compute cannot compensate for missing task-relevant signal.

B.2.4 Logits and numerical stability

Neural-network classifiers commonly output a vector of logits⁵ $\mathbf{z}$ (unnormalized scores). For class index i , softmax normalizes z_i over all class indices j to produce a probability:

$$\text{Softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

The problem is that if z_i is large (for example, 100), the exponential e^{z_i} overflows common training and inference formats such as FP32, BF16, and FP16. FP64 can represent this specific value, but production ML kernels rarely use FP64 for softmax. The solution is to compute in log-space: the “Log-Sum-Exp” trick allows us to compute $\log(\sum e^{z_j})$ without ever calculating the massive exponentials directly, preserving numerical precision.⁶

A small worked example shows why this trick matters in practice, even for large but realistic logit values:

Napkin Math B.3: Log-sum-exp in action

Setup: A three-class classifier outputs logits $\mathbf{z} = [100, 101, 102]$.

Without the trick (naive softmax):

$$\exp(100) \approx 2.7 \times 10^{43}, \exp(101) \approx 7.3 \times 10^{43}, \exp(102) \approx 2.0 \times 10^{44}$$

These numbers are representable in FP64 but overflow FP32 ($\max \approx 3.4 \times 10^{38}$). With FP16 ($\max \approx 65,504$), even e^{12} overflows. Raw logits rarely reach magnitude 100, but large intermediate values do arise (for example, unscaled attention scores or overconfident outputs), and because FP16 and BF16 are the standard training precisions, a naive softmax risks overflow often enough that the stable form is used by default.

With the trick: Subtract $a = \max(\mathbf{z}) = 102$:

Shifted logits: $\mathbf{z} - a = [-2, -1, 0]$

Exponentials: $\exp(-2) \approx 0.135$, $\exp(-1) \approx 0.368$, $\exp(0) = 1.0$

Sum = 1.503. LogSumExp = $102 + \log(1.503) = 102.408$.

Softmax: $[0.135/1.503, 0.368/1.503, 1.0/1.503] = [0.090, 0.245, 0.665]$

Systems insight: The exponentiated shifted values and the resulting softmax probabilities are in $[0, 1]$ —no overflow risk, even in FP16.

⁴ | **Entropy:** Shannon adopted the term in 1948 to quantify information content (Shannon 1948). In systems terms, Shannon entropy is a lower bound on average lossless code length. Practical prefix codes approach it, and block codes can approach it asymptotically. With $\log_2$, the unit is bits; natural logs give nats. These bounds inform compression ratios and data-pipeline sizing.

⁵ | **Logit:** Joseph Berkson coined the term in 1944 by analogy with *probit*. The logit function is the inverse of the logistic (sigmoid) function: $\text{logit}(p) = \log(p/(1-p))$. In deep learning, “logits” refers more loosely to the raw, unnormalized output of the final linear layer before any activation function is applied.

⁶ | **Log-sum-exp:** Implemented as `torch.logsumexp` in PyTorch and `scipy.special.logsumexp` in SciPy. It relies on the identity $\log(\sum e^{x_i}) = a + \log(\sum e^{x_i - a})$, where $a = \max(x_i)$. Shifted values $x_i - a$ are ≤ 0 , ensuring exponentials never overflow.

Together, these worked examples connect the full chain, from physical data movement to distribution shift and numerical stability.

B.3 Summary

The tools in this appendix—tail-aware metrics, drift divergences, information-theoretic bounds, and numerical stability techniques—support quantitative diagnosis of data-related failures rather than anecdotal explanations.

Algorithm Foundations

A profile showing low utilization is often misdiagnosed as a hardware limit when the cause is algorithmic: a matrix multiply shape that wastes tensor cores, a layout that breaks contiguity, or activation memory held for backpropagation. This appendix collects the computational complexity models and memory footprint derivations that the book’s performance chapters lean on, enabling readers to evaluate algorithmic trade-offs before committing code to accelerators. It assumes familiarity with linear algebra and with the iron law introduced early in the book.

How to Use This Appendix

This appendix is designed as a reference. Reach for it when translating a profiler symptom (“slow matmul,” “shape mismatch,” “OOM during training”) into a concrete computational or memory cause.

Conventions used here follow the book-wide notation (for example, B is reserved for batch size and BW denotes bandwidth).

- **When general matrix multiply (GEMM) kernels are slow:** Use section C.1.4 and compare intensity to the hardware’s ridge point.
- **When memory blows up in training:** Use section C.3.3 and the training memory decomposition.
- **When tensor code “should work” but does not:** Use section C.2.2 and section C.2.3.
- **When sparsity is proposed as a fix:** Use section C.1.5 to check density and metadata overhead.
- **When proving gradient compute bounds:** Use section C.4 for the Baur–Strassen reverse-mode automatic differentiation (AD) complexity theorem.
- **When analyzing sequence length bottlenecks:** Use section C.5 for the online softmax recurrence and FlashAttention high-bandwidth memory (HBM) I/O reduction proof.
- **When allocating compute between parameters and tokens:** Use section C.6 for the Kaplan and Chinchilla compute-optimal scaling law derivation.

Neural networks turn linear algebra into learned behavior: matrices transform activations, tensor layouts determine how data move through memory, and backpropagation carries error signals through the resulting computation. These foundations support the deep learning treatment in chapter 5, the framework internals in chapter 7, and the training strategies in chapter 8. Architectures change, but the underlying mathematical machinery remains constant.

C.1 Linear Algebra

Deep learning systems are, at their core, engines for transforming massive matrices. Frameworks like PyTorch abstract away the raw math, but performance engineering still depends on the underlying linear algebra. Building on how numbers are stored in section D.4, this section focuses on how they are manipulated.

Systems Perspective C.1: Why this matters

Many dense neural networks, especially transformers and large convolutional neural networks, spend most compute in matrix multiplication or GEMM-like kernels. A self-attention block

performs the Q, K, V, and output projection GEMMs plus two attention matrix multiplications: one for the attention scores and one for the attention-weighted values. The layer’s feed-forward block adds more GEMMs. Understanding GEMM performance characteristics explains why batch size affects throughput, why certain layer dimensions are “better” than others, and how to interpret profiler output. Reasoning about matrix dimensions and arithmetic intensity predicts whether a dense workload is compute bound or memory bound *before* any profiler trace runs.

C.1.1 Tensor operations and notation

Einstein summation¹ notation makes complex operations explicit throughout this book (implemented as `torch.einsum` in PyTorch and `np.einsum` in NumPy). Matrix multiplication $\mathbf{C} = \mathbf{AB}$ becomes:

$$C_{ij} = \sum_k A_{ik} B_{kj}$$

In einsum notation, this is `ik,kj->ij`. The notation extends naturally to the multi-dimensional operations in attention mechanisms. For example, batched multi-head attention is `bhid,bhjd->bhi j` (batch b , head h , query sequence i , key sequence j , and head dimension d).

C.1.2 Memory layouts and performance

Data layout in memory (row-major vs. column-major) directly affects cache efficiency. When iterating over a matrix, accessing contiguous memory locations is dramatically faster than strided access.

A common optimization pattern is to materialize a transposed tensor once before repeated operations to ensure contiguous access in the hot loop. The one-time copy cost is amortized across many subsequent operations.

C.1.3 The dot product as similarity

The dot product $\mathbf{a} \cdot \mathbf{b} = \sum a_i b_i$ is geometrically equivalent to $\|\mathbf{a}\| \|\mathbf{b}\| \cos \phi$, which makes it a natural measure of similarity between two vectors: for nonzero vectors, its sign distinguishes acute, right, and obtuse angles, while its magnitude also depends on both vector norms.

This geometric interpretation is why dot products appear everywhere in modern architectures. In attention mechanisms, query (Q) and key (K) vectors are dot-produced to compute a similarity score that determines how much each token attends to every other token. The resulting attention weights are then used to form a weighted combination of value (V) vectors—making the dot product the foundation of the transformer’s ability to model long-range dependencies.

C.1.4 General matrix multiply (GEMM)

GEMM² is the computational workhorse of deep learning. For matrices of size $M \times K$ and $K \times N$, GEMM performs $2MNK$ floating-point operations (multiply-accumulate counts as two operations).

The arithmetic intensity of GEMM scales linearly with matrix dimension. For square $n \times n$ matrices in FP16 (2 bytes/element), the ideal $\beta = 0$ bound reads $\mathbf{A}$ and $\mathbf{B}$ once and writes $\mathbf{C}$ once:

$$\text{Intensity} = \frac{O}{D_{\text{vol}}} = \frac{2n^3}{3n^2 \times 2} = \frac{n}{3} \text{ FLOP/byte}$$

Reading the old $\mathbf{C}$ when $\beta \neq 0$ changes the intensity to $n/4$ FLOP/byte. This explains several important phenomena:

- **Larger batches can improve efficiency:** Batching increases effective matrix dimensions when it forms larger GEMMs or enables reuse; merely queuing independent small GEMMs does not raise intensity.
- **Aligned dimensions help:** Hardware tensor cores are optimized for precision- and architecture-specific tile multiples. Dimensions that align with these multiples avoid padding overhead and improve kernel efficiency, but they do not need to be powers of two.
- **Small matrices are inefficient:** A square GEMM with $n = 64$ has intensity $64/3 \approx 21.3$ FLOP/byte, well below the ridge point (153.0 FLOP/byte). Its roofline cap is ~ 13.9 percent of peak; launch and tiling overhead can lower realized throughput.

¹ | **Einstein summation convention:** Repeated indices in a product are implicitly summed over, eliminating explicit summation signs. ML frameworks adopted the convention because it concisely expresses arbitrary tensor contractions in a single string.

² | **General matrix multiply (GEMM):** The BLAS (Basic Linear Algebra Subprograms) family began with vector routines standardized in 1979 (Lawson et al. 1979); GEMM was standardized later as a Level 3 routine (Dongarra et al. 1988). The “GE” prefix stands for “general,” as opposed to symmetric or triangular forms. GEMM computes $\mathbf{C} = \alpha \mathbf{AB} + \beta \mathbf{C}$ and is a performance-critical routine in deep learning. Section 8.3.1.2 explains how GEMM shapes determine training throughput.

C.1.5 Sparse matrix formats

When most elements in a matrix are zero, specialized storage formats avoid wasting memory on zeros and enable computations that skip them entirely. The **compressed sparse row (CSR)** format uses three arrays:

- **Values:** The nonzero elements, stored in row order
- **Col_Idx:** The column index of each nonzero element
- **Row_Ptr:** The starting position in **Values** for each row (length = num_rows + 1)

CSR is useful for sparse feature matrices in recommendation pipelines and for pruned model weights. For a matrix with N elements, R rows, and K nonzeros, CSR uses $\mathcal{O}(K + R)$ storage instead of $\mathcal{O}(N)$; when K is large relative to R , this is often summarized as $\mathcal{O}(K)$.

To see the trade-off concretely, consider a vocabulary embedding matrix with 100,000 rows and 10,000 columns (1B parameters):

- **Dense (FP32):** 1B parameters $\times$ 4 bytes each = 4 GB.
- **Sparse (1 percent density):** CSR stores roughly 10M entries $\times$ (4-byte value + 4-byte column index), plus one row pointer per row and a final pointer, totaling $\approx$ 80.4 MB.
- **Result:** A 50 $\times$ reduction in memory footprint, fitting a model that would otherwise OOM (Out of Memory).

Linear algebra tells us *what* to compute; the next question is *how* to express those computations in code. Tensor programming primitives—shapes, strides, and broadcasting—bridge the gap between mathematical notation and the array operations that actually execute on hardware.

C.2 Tensor Programming Primitives

A shape mismatch crash, a silently wrong broadcast, a kernel running at 5 percent of peak because of a noncontiguous tensor—these common ML engineering failures all trace back to the same layer of abstraction. *Tensor programming* translates the abstract math of linear algebra into concrete array manipulations that run on hardware.

C.2.1 Computational complexity cheat sheet

Table C.1 provides a quantitative reference for the most common building blocks. Use these formulas for napkin-math estimation of model size and compute requirements *before* hardware is provisioned. Given layer dimensions and input shapes, the formulas provide first-order parameter and compute estimates. The attention row is the key warning: sequence length contributes an S^2 term, while hidden dimension contributes d^2 projection work, so long-context models can shift the dominant cost without changing the parameter count.

Table C.1: Deep Learning Tensor Primitives: Summary of shapes, parameters, and FLOP counts. Note: B is batch size, S is sequence length, and K is kernel size. The attention FLOPs include QKV projections and score interactions. The S^2 term applies to full-sequence training and prefill, while KV-cached generation is linear per token in cached context; d^2 projection work dominates at large hidden dimensions.

Layer Type	Output Shape	Parameters (P)	FLOPs (per Forward Pass)
Linear	(B, N_{out})	$N_{\text{in}} \times N_{\text{out}} + N_{\text{out}}$ if bias is enabled	$2 \times B \times N_{\text{in}} \times N_{\text{out}}$
Conv2D	$(B, C_{\text{out}}, H', W')$	$K^2 \times C_{\text{in}} \times C_{\text{out}} + C_{\text{out}}$ if bias is enabled	$2 \times B \times H' \times W' \times K^2 \times C_{\text{in}} \times C_{\text{out}}$
Multi-Head Self-Attention	(B, S, d_{model})	$4 \times d_{\text{model}}^2$, plus $4 \times d_{\text{model}}$ if projection biases are enabled	$B \times (4S^2 d_{\text{model}} + 8Sd_{\text{model}}^2)$
LayerNorm	(B, S, d_{model})	$2 \times d_{\text{model}}$ if affine scale and bias are enabled	$\mathcal{O}(B \times S \times d_{\text{model}})$

C.2.2 Shapes and strides

A tensor is a view over an underlying storage buffer, described by shape, stride, dtype, and offset metadata. Only contiguous tensors lay their logical elements out as one adjacent block.

- **Shape:** The dimensions of the tensor (for example, (3, 4)).
- **Stride:** The number of elements to skip in memory to move to the next element in a dimension.


Operations like `transpose()` or `view()` often just change the *strides*, not the data in memory. This is fast ($\mathcal{O}(1)$) but can lead to noncontiguous tensors that fail in kernels requiring contiguous data. In such cases, calling `contiguous()` forces an $\mathcal{O}(N)$ memory copy that can dominate runtime if triggered repeatedly inside a loop.

C.2.3 Broadcasting

Broadcasting allows arithmetic operations on tensors of different shapes. Compare dimensions from the last to the first. Two dimensions are compatible when:

1. They are equal.
2. One of them is 1.

The dimension with size one is “stretched” to match the other, as illustrated in figure C.1. This stretching is virtual: the data is not copied in memory. Instead, the stride for that dimension is set to 0, allowing the hardware to read the same value repeatedly with $\mathcal{O}(1)$ memory overhead.



Dashed cells are virtually expanded via stride=0 (no memory allocation)

Figure C.1: Tensor Broadcasting Rules: A (3,1) column and a (1,4) row broadcast across complementary dimensions, producing the pairwise sums in a (3,4) result.

Consider a concrete case: tensor A has shape (32, 1, 64) and tensor B has shape (1, 128, 64). Comparing dimensions right to left, 64 matches 64, then 1 stretches to 128, then 1 stretches to 32, yielding result shape (32, 128, 64). Visualizing this expansion prevents silent logic bugs that accidentally allocate a large tensor (for example, a (Batch, Batch) matrix instead of an element-wise (Batch) vector).

Shapes, strides, and broadcasting govern how tensors flow through a model’s forward pass. Training adds a second requirement: *learning from errors*. Backpropagation makes that learning possible and imposes the memory costs developed below.

C.3 Mechanics of Learning

Valid tensor programs make training loops possible. Backpropagation orchestrates these tensors to compute gradients, transforming a forward prediction into a backward learning signal.

🔍 Systems Perspective C.2: Why this matters

When training fails—loss goes to NaN, gradients explode, or memory runs out—understanding what backpropagation actually does is essential for diagnosing the problem. The backward graph supplies the mental model for reasoning about gradient flow and memory usage during training.

C.3.1 The chain rule and automatic differentiation

For a composed function $y = f(g(x))$, the derivative is $\frac{dy}{dx} = \frac{dy}{dg} \cdot \frac{dg}{dx}$. In a neural network, f and g are layers, and the composition can be many levels deep. For a three-layer network $y = f_3(f_2(f_1(x)))$, the chain rule extends to:

$$\frac{\partial y}{\partial x} = \frac{\partial f_3}{\partial f_2} \cdot \frac{\partial f_2}{\partial f_1} \cdot \frac{\partial f_1}{\partial x}$$

Each factor in this product is a *local* derivative—computed at one layer using only that layer’s inputs and outputs. This locality is what makes the algorithm tractable: the entire network never needs to be differentiated as a monolithic function. Instead, each layer computes its own local derivative during the backward pass and multiplies it by the gradient flowing in from the layer above.

Modern frameworks use reverse-mode automatic differentiation, which computes gradients for all P parameters in a single backward pass. The key insight is that starting from the output and working backward (reverse mode) requires one pass regardless of the number of parameters, whereas starting from each input and working forward (forward mode) would require P passes—one per parameter. This is why a training step is a small constant multiple of inference, commonly about $2\text{--}3\times$ a forward pass for dense networks, rather than P passes.

C.3.2 The backpropagation algorithm

Backpropagation³ implements the chain rule efficiently through two passes: forward to compute outputs, backward to compute gradients. Figure C.2 illustrates this process for a simple two-layer network, with the forward pass (gray arrows) computing outputs and the backward pass (red dashed arrows) propagating gradients.




Figure C.2: Backpropagation Computational Graph: Solid gray arrows carry the forward computation from input x to loss $\mathcal{L}$, while dashed red arrows carry gradients backward along the same dependencies.

The gradient edges retrace the forward dependencies, so each backward step requires values retained from the corresponding forward operation.

Forward pass

Using row-vector activations, omitting biases, and retaining intermediate activations for backward, start at x , the input. Multiply by W_1 to get hidden activation h . Cache h because the backward pass will need it later. Multiply h by W_2 to get output y . Cache y . Compare y to the target label to compute loss $\mathcal{L}$.

At this point, the loss has been computed and memory contains the input x , the cached activation h , the cached output y , and the loss $\mathcal{L}$. For a large model, these cached activations can dominate memory usage.

Backward pass

Now trace backward from $\mathcal{L}$. The loss function provides $\frac{\partial \mathcal{L}}{\partial y}$, the gradient of loss with respect to the prediction. This is where the error signal enters the network.

Since $y = h \cdot W_2$, the chain rule gives us two gradients at this layer:

$$\frac{\partial \mathcal{L}}{\partial W_2} = h^T \cdot \frac{\partial \mathcal{L}}{\partial y} \quad (\text{weight gradient—used to update } W_2)$$

$$\frac{\partial \mathcal{L}}{\partial h} = \frac{\partial \mathcal{L}}{\partial y} \cdot W_2^T \quad (\text{input gradient—passed backward to the previous layer})$$

³ | **Backpropagation:** Short for “backward propagation of errors.” The algorithm was independently discovered multiple times—by Werbos (1974) and Linnainmaa (1970) for reverse-mode differentiation and by Rumelhart et al. (1986) for neural network training. Its key insight is that computing gradients for *all* parameters requires one backward pass through the graph rather than one pass per parameter. Section C.4 derives the constant-factor work bound and a representative dense-layer cost.

Computing $\frac{\partial \mathcal{L}}{\partial W_2}$ requires the cached activation h from the forward pass. Activations must remain available during backward because every layer's weight gradient depends on that layer's input.

Now continue backward to the first layer. Since $h = x \cdot W_1$, the same pattern gives:

$$\frac{\partial \mathcal{L}}{\partial W_1} = x^T \cdot \frac{\partial \mathcal{L}}{\partial h}$$

Each step backward requires two things: the gradient flowing in from above ($\frac{\partial \mathcal{L}}{\partial h}$) and that layer's input (x), which must be available during backward. This is why the backward pass costs roughly $2\times$ the forward pass in compute: at each layer, it performs two matrix multiplications (one for the weight gradient, one for the input gradient) vs. one in the forward pass.

C.3.3 The true cost of training memory

A common mistake is to assume that training memory equals model size. This assumption leads to immediate OOM errors because weights are only one of four components. Training memory comprises weights, gradients, optimizer state, and retained activations:

$$M_{\text{total}} = M_{\text{weights}} + M_{\text{gradients}} + M_{\text{optimizer}} + M_{\text{activations}}$$

For a standard Adaptive Moment Estimation (Adam) optimizer (Kingma and Ba 2015) in mixed precision (Micikevicius et al. 2017):

- **Weights:** 2 bytes (FP16/BF16) or 4 bytes (FP32).
- **Gradients:** Same size as weights.
- **Optimizer State:** Adam moments require 8 bytes per parameter; an FP32 master copy adds another 4 bytes, giving 8–12 bytes per parameter.
- **Activations:** The hidden giant. A recomputation- or tiled-attention-friendly estimate scales as $\mathcal{O}(B \times S \times N_L \times d)$; a full no-recompute attention implementation also materializes attention-score tensors that scale with S^2 .

To see how these components interact in practice, consider a concrete model.



Napkin Math C.1: Worked example: GPT-2 (1.5B) training memory

The model: GPT-2 XL has $P = 1.5 \times 10^9$ parameters, 48 layers, hidden dimension $d = 1600$.

Model state (fixed per step):

- **Weights (BF16):** $1.5 \times 10^9 \times 2 \text{ bytes} = 3 \text{ GB}$
- **Gradients (BF16):** $1.5 \times 10^9 \times 2 \text{ bytes} = 3 \text{ GB}$
- **Optimizer (FP32 master + momentum + variance):** $1.5 \times 10^9 \times 12 \text{ bytes} = 18 \text{ GB}$
- **Total model state:** 24 GB—fits on an 80 GB-class A100/H100 (85.9 GB in decimal units), but leaves only 61.9 GB for activations.

Activations (scale with batch):

Per-layer retained activations for this quick estimate are approximately $12 \times B \times S \times d$ BF16 elements, or $12 \times B \times S \times d \times 2$ bytes, where B is batch size and S is sequence length. The factor twelve accounts for the major intermediate tensors retained for backpropagation: input activations, QKV projections ($3d$), attention output, FFN intermediate ($4d$), and layer norm/dropout masks. This estimate assumes the implementation does not retain the full $B \times N_{\text{heads}} \times S^2$ attention-score tensor, as in recomputation- or tiled-attention-friendly accounting. With 48 layers, batch size 8, and sequence length 1024:

$$48 \times 12 \times 8 \times 1024 \times 1600 \times 2 \text{ bytes} \approx 15.1 \text{ GB}$$

Systems insight: ~39.1 GB under this activation-accounting convention, which fits on one 85.9 GB accelerator. However, increase the batch to 64 and activations grow to ~120.8 GB, exceeding the remaining capacity. This is the threshold where gradient checkpointing or

another memory-saving method becomes necessary; its savings and recomputation overhead depend on checkpoint placement and implementation.

C.3.3.1 Activation explosion

While weights are fixed ($\mathcal{O}(P)$), activations grow linearly with batch size and at least linearly with sequence length; implementations that materialize attention scores add an $\mathcal{O}(S^2)$ component. As the worked example shows, activation memory can quickly exceed the room left after model state. Gradient checkpointing⁴ reduces this pressure by storing only selected activations and recomputing the rest during backpropagation (T. Chen et al. 2016); techniques such as FlashAttention address related attention-memory bottlenecks by tiling attention to reduce memory round-trips (T. Dao et al. 2022).

The same decomposition provides a quick feasibility check: sum the per-parameter model state, add the activation term implied by batch size and sequence length, and compare the total against the accelerator’s capacity.

C.3.4 Computational graphs and optimization

The dependency structure exposed by backpropagation also gives compilers something to optimize. ML compilers represent models as directed acyclic graphs (DAGs), and that representation enables hardware-independent transformations.

C.3.4.1 Static single assignment

Compilers transform graphs into static single-assignment (SSA) form, where each variable is assigned exactly once. This makes data dependencies explicit, enabling safe optimizations—most importantly, *operator fusion*.

C.3.4.2 Operator fusion

Without fusion, each operation in a chain like $\text{MatMul} \rightarrow \text{Add (bias)} \rightarrow \text{ReLU}$ produces an intermediate tensor that is written to HBM and then read back for the next operation. For elementwise operations like Add and ReLU, the compute is trivial (one FLOP per element) but the memory traffic is not (read the tensor, write it back). The arithmetic intensity of unfused elementwise operations is therefore close to zero—deeply memory bound.

Fusion combines consecutive operations into a single kernel that reads the input once, applies all operations in registers or shared memory, and writes the final result once. For a sequence of k elementwise operations on a tensor of size N bytes, fusion reduces memory traffic from $2kN$ bytes (each op reads and writes) to $2N$ bytes (one read, one write)—a $k\times$ reduction.

FlashAttention is an I/O-aware tiled attention algorithm: it computes the score, softmax, and value-product stages in SRAM tiles without materializing the full attention matrix in HBM, reducing HBM traffic and auxiliary attention memory from $\mathcal{O}(S^2)$ to $\mathcal{O}(S)$. The original work reports up to $3\times$ wall-clock speedups on its evaluated workloads (T. Dao et al. 2022); realized gains depend on shape and hardware (see section C.5 for the step-by-step derivation). Together, the linear algebra foundations, tensor programming mechanics, and training memory model covered in this appendix form the algorithmic substrate on which all ML systems are built.

C.4 Baur–Strassen Reverse-Mode AD Complexity Theorem

Modern deep learning scales to hundreds of billions of parameters because evaluating the full gradient vector $\nabla f(x)$ of a scalar loss function $f : \mathbb{R}^N \rightarrow \mathbb{R}$ requires compute proportional to evaluating $f(x)$ itself, with a constant-factor bound independent of the input dimension N . Baur and Strassen proved this complexity result in 1983 (Baur and Strassen 1983).

C.4.1 Mathematical formulation and theorem statement

Let $f : \mathbb{R}^N \rightarrow \mathbb{R}$ be a multivariate scalar-valued function represented as a computational graph $G = (V, E)$ consisting of $W(f)$ elementary operations (additions, subtractions, multiplications, divisions, and unary functions such as $\exp, \ln, \sin, \sqrt{\cdot}$).

⁴ | **Gradient checkpointing:** Trades compute for memory. Instead of storing all activations, only a subset (checkpoints) is kept, and the missing ones are recomputed during the backward pass. This reduces memory usage from storing every layer activation to roughly the square root of the layer count at the cost of ~ 33 percent more compute (T. Chen et al. 2016).


Theorem C.1: Baur–Strassen complexity theorem

If a scalar function $f : \mathbb{R}^N \rightarrow \mathbb{R}$ can be computed in $W(f)$ elementary operations, then the full gradient vector $\nabla f(x) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_N} \right)^T \in \mathbb{R}^N$ can be evaluated in total work $\text{Work}_{\text{Reverse}}(f, \nabla f)$ satisfying:

$$\text{Work}_{\text{Reverse}}(f, \nabla f) \leq C \cdot W(f)$$

where $C \leq 5$ is a universal constant independent of the input dimension N . Furthermore, the backward pass alone requires work $\text{Work}_{\text{Backward}}(\nabla f) \leq 4 \cdot W(f)$.

C.4.2 Step-by-step chain rule adjoint graph propagation proof

Let the computational graph $G = (V, E)$ be topologically ordered with nodes $v_{1-N}, \dots, v_0, v_1, v_2, \dots, v_W$, where input variables are $x_i = v_{i-N}$ for $i \in \{1, \dots, N\}$ and the final output is $f(x) = v_W$. Each intermediate node v_j is evaluated via a primitive operation:

$$v_j = \phi_j(\{v_k \mid k \in \text{Parents}(j)\})$$

Define the *adjoint variable* $\bar{v}_i$ for every node v_i as the partial derivative of the scalar output v_W with respect to v_i :

$$\bar{v}_i \triangleq \frac{\partial f}{\partial v_i} = \frac{\partial v_W}{\partial v_i}$$

By the multivariate chain rule, the adjoint $\bar{v}_i$ is the sum of partial derivatives propagated from all successor nodes (children) that consume v_i :

$$\bar{v}_i = \sum_{j \in \text{Children}(i)} \bar{v}_j \cdot \frac{\partial v_j}{\partial v_i}$$

Reverse-mode automatic differentiation initializes $\bar{v}_W = 1$ and $\bar{v}_i = 0$ for all $i < W$, then traverses the nodes in reverse topological order from v_W down to v_1 . For each primitive node v_j , its fully accumulated adjoint $\bar{v}_j$ propagates error signals to its parent nodes v_k ($k \in \text{Parents}(j)$) via:

$$\bar{v}_k \leftarrow \bar{v}_k + \bar{v}_j \cdot \frac{\partial v_j}{\partial v_k}$$

We now perform operation-by-operation micro-cost accounting for all primitive nodes to bound the backward work:

1. Addition/Subtraction ($v_j = v_i \pm v_k$):

- Local derivatives: $\frac{\partial v_j}{\partial v_i} = 1, \frac{\partial v_j}{\partial v_k} = \pm 1$.
- Backward updates: $\bar{v}_i \leftarrow \bar{v}_i + \bar{v}_j$ and $\bar{v}_k \leftarrow \bar{v}_k \pm \bar{v}_j$.
- Cost: 2 additions in backward pass for 1 forward operation. Work ratio: $\text{Fwd}(1) + \text{Bwd}(2) = 3 \text{ OPs} \leq 5 \cdot 1$.

2. Multiplication ($v_j = v_i \cdot v_k$):

- Local derivatives: $\frac{\partial v_j}{\partial v_i} = v_k, \frac{\partial v_j}{\partial v_k} = v_i$.
- Backward updates: $\bar{v}_i \leftarrow \bar{v}_i + \bar{v}_j \cdot v_k$ and $\bar{v}_k \leftarrow \bar{v}_k + \bar{v}_j \cdot v_i$.
- Cost: 2 multiplications + 2 additions = 4 operations in backward pass for 1 forward operation. Work ratio: $\text{Fwd}(1) + \text{Bwd}(4) = 5 \text{ OPs} \leq 5 \cdot 1$.

3. Division ($v_j = v_i / v_k$):

- Local derivatives: $\frac{\partial v_j}{\partial v_i} = \frac{1}{v_k}, \frac{\partial v_j}{\partial v_k} = -\frac{v_i}{v_k^2} = -\frac{v_j}{v_k}$.
- Backward updates: $\bar{v}_i \leftarrow \bar{v}_i + \frac{\bar{v}_j}{v_k}$ and $\bar{v}_k \leftarrow \bar{v}_k - \frac{\bar{v}_j \cdot v_j}{v_k}$.

- Cost: 2 multiplications/divisions + 2 additions = 4 operations in backward pass for 1 forward operation. Work ratio: $\text{Fwd}(1) + \text{Bwd}(4) = 5 \text{ OPs} \leq 5 \cdot 1$.

4. Unary Transcendental Operations ($v_j = \phi(v_i)$):

- Local derivative: $\frac{\partial v_i}{\partial v_i} = \phi'(v_i)$.
- Backward update: $\bar{v}_i \leftarrow \bar{v}_i + \bar{v}_j \cdot \phi'(v_i)$.
- For $v_j = \exp(v_i)$, $\phi'(v_i) = v_j$, requiring 1 multiplication + 1 addition = 2 OPs.
- For $v_j = \ln(v_i)$, $\phi'(v_i) = 1/v_i$, requiring 1 division + 1 addition = 2 OPs.

Summing across all primitive nodes in G , every forward operation generates at most 4 backward operations:

$$\text{Work}_{\text{Backward}}(\nabla f) \leq 4 \cdot W(f)$$

Adding the forward pass work $W(f)$ yields the total work bound:

$$\text{Work}_{\text{Reverse}}(f, \nabla f) = W(f) + \text{Work}_{\text{Backward}}(\nabla f) \leq 5 \cdot W(f) \quad \blacksquare$$

Systems Perspective C.3: Why 175B-parameter large language models can be trained

The independence of $C \leq 5$ from input dimension N is the mathematical reason deep learning works. Consider training a transformer model with $P = 175 \times 10^9$ parameters:

- **Forward Pass:** Matrix multiplications $Y = XW$ require $2P$ FLOPs per token.
- **Reverse-Mode Backward Pass:** Computes both weight gradients $\frac{\partial \mathcal{L}}{\partial W} = X^T \frac{\partial \mathcal{L}}{\partial Y}$ ($2P$ FLOPs) and input activation gradients $\frac{\partial \mathcal{L}}{\partial X} = \frac{\partial \mathcal{L}}{\partial Y} W^T$ ($2P$ FLOPs), totaling $4P$ FLOPs per token.
- **Total Training Work:** $2P(\text{fwd}) + 4P(\text{bwd}) = 6P$ FLOPs per token. The backward compute factor is exactly $\frac{4P}{2P} = 2$, well within the Baur–Strassen theoretical bound of $C \leq 5$.

If we instead used forward-mode AD or finite differences to compute $\nabla f \in \mathbb{R}^P$, evaluating each partial derivative separately would require P forward passes: $\text{Work}_{\text{ForwardMode}}(\nabla f) = \mathcal{O}(P \cdot W) = \mathcal{O}(P^2)$ FLOPs. For $P = 175 \times 10^9$, that method would require $2 \times (175 \times 10^9)^2 \approx 6.1 \times 10^{22}$ FLOPs per token— P times the forward-pass work and therefore prohibitive at this scale. Reverse-mode AD collapses this $\mathcal{O}(P)$ factor into a constant multiplier of 2 for the dense-layer accounting above.

C.5 FlashAttention Tiling & Online Softmax Recurrence Proof

Standard self-attention in transformers requires materializing an $N \times N$ attention score matrix in HBM, creating an $\mathcal{O}(N^2)$ memory bandwidth bottleneck. FlashAttention (T. Dao et al. 2022) eliminates this bottleneck by tiling query, key, and value matrices and recomputing attention using an online softmax recurrence without materializing intermediate scores to HBM.

C.5.1 Standard softmax HBM I/O bottleneck

Given queries $\mathbf{Q} \in \mathbb{R}^{N \times d}$, keys $\mathbf{K} \in \mathbb{R}^{N \times d}$, and values $\mathbf{V} \in \mathbb{R}^{N \times d}$ for sequence length N and head dimension d , standard attention evaluates:

$$\mathbf{S} = \mathbf{QK}^T \in \mathbb{R}^{N \times N}, \quad \mathbf{A} = \text{softmax}(\mathbf{S}) \in \mathbb{R}^{N \times N}, \quad \mathbf{O} = \mathbf{AV} \in \mathbb{R}^{N \times d}$$

To prevent numerical overflow, row-wise softmax subtracts the row maximum $m_i = \max_{1 \leq j \leq N} S_{ij}$:

$$A_{ij} = \frac{\exp(S_{ij} - m_i)}{\sum_{k=1}^N \exp(S_{ik} - m_i)}$$

Standard GPU implementations execute this in three discrete kernel launches:

1. Compute $\mathbf{S} = \mathbf{QK}^T$ in SRAM, then write $\mathbf{S} \in \mathbb{R}^{N \times N}$ to HBM ($\Theta(N^2)$ writes).
2. Read $\mathbf{S}$ from HBM, compute row max m and row sum d , compute $\mathbf{A} = \text{softmax}(\mathbf{S})$, write $\mathbf{A} \in \mathbb{R}^{N \times N}$ to HBM ($\Theta(N^2)$ reads + writes).
3. Read $\mathbf{A}$ and $\mathbf{V}$ from HBM, compute $\mathbf{O} = \mathbf{AV}$, write $\mathbf{O} \in \mathbb{R}^{N \times d}$ to HBM ($\Theta(N^2 + Nd)$ reads + writes).

Total HBM memory traffic is $\Theta(N^2 + Nd)$ elements. For $N = 8192$, $d = 128$ in FP16, $\mathbf{S}$ requires $8192^2 \times 2 \text{ B} = 134 \text{ MB}$ per head per layer. Because modern GPUs feature massive compute throughput relative to HBM bandwidth, repeatedly reading and writing intermediate matrices $\mathbf{S}$ and $\mathbf{A}$ leaves GPU execution units memory-bound.

C.5.2 Online softmax recurrence proof

Tiling softmax is difficult because computing $d_i = \sum_{k=1}^N \exp(S_{ik} - m_i)$ requires the global maximum m_i across all N keys. The online softmax recurrence solves this by incrementally updating normalization statistics as sub-blocks arrive.



Theorem C.2: Online softmax recurrence theorem

Let a vector $x \in \mathbb{R}^N$ be partitioned into two sub-vectors $x^{(1)} \in \mathbb{R}^{N_1}$ and $x^{(2)} \in \mathbb{R}^{N_2}$ with $N = N_1 + N_2$. Let the local statistics for block 1 be:

$$m^{(1)} = \max_{1 \leq j \leq N_1} x_j^{(1)}, \quad d^{(1)} = \sum_{j=1}^{N_1} \exp(x_j^{(1)} - m^{(1)})$$

and for block 2 be:

$$m^{(2)} = \max_{1 \leq j \leq N_2} x_j^{(2)}, \quad d^{(2)} = \sum_{j=1}^{N_2} \exp(x_j^{(2)} - m^{(2)})$$

Then the combined maximum m^{new} , normalization sum d^{new} , and accumulated output vector $\mathbf{O}^{\text{new}}$ satisfy the exact recurrence relations:

$$\begin{aligned} m^{\text{new}} &= \max(m^{(1)}, m^{(2)}) \\ d^{\text{new}} &= d^{(1)} e^{m^{(1)} - m^{\text{new}}} + d^{(2)} e^{m^{(2)} - m^{\text{new}}} \\ \mathbf{O}^{\text{new}} &= \frac{d^{(1)} e^{m^{(1)} - m^{\text{new}}}}{d^{\text{new}}} \mathbf{O}^{(1)} + \frac{e^{m^{(2)} - m^{\text{new}}}}{d^{\text{new}}} \sum_{j=1}^{N_2} e^{x_j^{(2)} - m^{(2)}} \mathbf{V}_j^{(2)} \end{aligned}$$

where $\mathbf{O}^{(1)} = \frac{1}{d^{(1)}} \sum_{j=1}^{N_1} e^{x_j^{(1)} - m^{(1)}} \mathbf{V}_j^{(1)}$ is the partial output accumulated from block 1.

Derivation of denominator recurrence

Expanding the combined sum $d^{\text{new}} = \sum_{j=1}^N e^{x_j - m^{\text{new}}}$:

$$\begin{aligned} d^{\text{new}} &= \sum_{j=1}^{N_1} e^{x_j^{(1)} - m^{\text{new}}} + \sum_{j=1}^{N_2} e^{x_j^{(2)} - m^{\text{new}}} \\ &= \sum_{j=1}^{N_1} e^{(x_j^{(1)} - m^{(1)}) + (m^{(1)} - m^{\text{new}})} + \sum_{j=1}^{N_2} e^{(x_j^{(2)} - m^{(2)}) + (m^{(2)} - m^{\text{new}})} \\ &= e^{m^{(1)} - m^{\text{new}}} \underbrace{\sum_{j=1}^{N_1} e^{x_j^{(1)} - m^{(1)}}}_{d^{(1)}} + e^{m^{(2)} - m^{\text{new}}} \underbrace{\sum_{j=1}^{N_2} e^{x_j^{(2)} - m^{(2)}}}_{d^{(2)}} \\ &= d^{(1)} e^{m^{(1)} - m^{\text{new}}} + d^{(2)} e^{m^{(2)} - m^{\text{new}}}. \quad \blacksquare \end{aligned}$$

Derivation of output rescaling recurrence

The true normalized output over the concatenated sequence is $\mathbf{O}^{\text{new}} = \frac{1}{d^{\text{new}}} \sum_{j=1}^N e^{x_j - m^{\text{new}}} \mathbf{V}_j$:

$$\begin{aligned} \mathbf{O}^{\text{new}} &= \frac{1}{d^{\text{new}}} \left[\sum_{j=1}^{N_1} e^{x_j^{(1)} - m^{\text{new}}} \mathbf{V}_j^{(1)} + \sum_{j=1}^{N_2} e^{x_j^{(2)} - m^{\text{new}}} \mathbf{V}_j^{(2)} \right] \\ &= \frac{1}{d^{\text{new}}} \left[e^{m^{(1)} - m^{\text{new}}} \sum_{j=1}^{N_1} e^{x_j^{(1)} - m^{(1)}} \mathbf{V}_j^{(1)} + e^{m^{(2)} - m^{\text{new}}} \sum_{j=1}^{N_2} e^{x_j^{(2)} - m^{(2)}} \mathbf{V}_j^{(2)} \right] \end{aligned}$$

Noting that $d^{(1)} \mathbf{O}^{(1)} = \sum_{j=1}^{N_1} e^{x_j^{(1)} - m^{(1)}} \mathbf{V}_j^{(1)}$, we substitute this directly:

$$\mathbf{O}^{\text{new}} = \frac{d^{(1)} e^{m^{(1)} - m^{\text{new}}}}{d^{\text{new}}} \mathbf{O}^{(1)} + \frac{e^{m^{(2)} - m^{\text{new}}}}{d^{\text{new}}} \sum_{j=1}^{N_2} e^{x_j^{(2)} - m^{(2)}} \mathbf{V}_j^{(2)}. \quad \blacksquare$$

C.5.3 Formal HBM access complexity proof

Let M be the fast SRAM capacity (in words). Following Algorithm 1 of T. Dao et al. (2022), FlashAttention sets $B_c = \lfloor M/(4d) \rfloor$ and $B_r = \min(\lfloor M/(4d) \rfloor, d)$. The proof abstracts constant factors and buffer reuse, using $B_c = \Theta(M/d)$, $B_r = \Theta(\min(M/d, d))$, and $B_r B_c = O(M)$ for the on-chip working set.

The algorithm executes two nested loops:

- Partition $\mathbf{Q}$ into $T_r = \lceil N/B_r \rceil$ blocks of size $B_r \times d$.
- Partition $\mathbf{K}, \mathbf{V}$ into $T_c = \lceil N/B_c \rceil$ blocks of size $B_c \times d$.
- **Outer loop** over $j = 1, \dots, T_c$: Load $\mathbf{K}_j$ and $\mathbf{V}_j$ into SRAM.
- **Inner loop** over $i = 1, \dots, T_r$: Load $\mathbf{Q}_i$, $\mathbf{O}_i$, and its softmax statistics; compute $\mathbf{S}_{ij} = \mathbf{Q}_i \mathbf{K}_j^T$ in SRAM; then update and write $\mathbf{O}_i$ and the statistics without materializing $\mathbf{S}_{ij}$ in HBM.

This loop order is the source of the I/O reduction. A key-value tile remains resident while the algorithm visits every query tile, so each element of $\mathbf{K}$ and $\mathbf{V}$ crosses the HBM boundary once. The query and output tiles must be revisited for each key-value block, but the $B_r \times B_c$ score tile is consumed on chip and discarded. Standard attention instead writes and rereads the full $N \times N$ score and probability matrices.

Total memory access accounting

1. **Key and value reads**: Each element of $\mathbf{K}$ and $\mathbf{V}$ is loaded once, contributing $2Nd$ accesses.
2. **Query, output, and statistics passes**: Each of the T_c column blocks reads all of $\mathbf{Q}$, reads and writes $\mathbf{O}$, and updates $O(N)$ softmax statistics. The matrices dominate the statistics for head dimension $d \geq 1$, so each pass contributes $\Theta(Nd)$ accesses and all passes contribute $\Theta(NdT_c)$.

Because $T_c = \lceil N/B_c \rceil = \Theta(Nd/M)$, the total is

$$\text{HBM}_{\text{Flash}} = \Theta(Nd + NdT_c) = \Theta\left(Nd + \frac{N^2 d^2}{M}\right) = \Theta\left(\frac{N^2 d^2}{M}\right),$$

where the final equality follows for $d \leq M \leq Nd$.

The bound counts communication rather than arithmetic. FlashAttention still evaluates the same dense score blocks and may recompute intermediate quantities during backpropagation; its advantage comes from keeping those intermediates out of HBM.

I/O reduction ratio

Comparing standard attention I/O to FlashAttention I/O:

$$\frac{\text{HBM}_{\text{Standard}}}{\text{HBM}_{\text{Flash}}} = \frac{\Theta(N^2 + Nd)}{\Theta\left(Nd + \frac{N^2 d^2}{M}\right)} = \Theta\left(\frac{M}{d^2}\right) \quad \text{for } N \gg \max\left(d, \frac{M}{d}\right)$$

🔍 Systems Perspective C.4: A100 I/O reduction

The asymptotic ratio M/d^2 predicts how greater SRAM capacity can reduce HBM traffic, but hidden constants and kernel details prevent it from serving as an exact transfer count. In an NVIDIA A100 measurement, T. Dao et al. (2022) report the effect for a forward-and-backward attention workload with sequence length 1,024, head dimension 64, 16 heads, batch size 64, a key-padding mask, and no dropout:

- **Standard attention:** 35.3 GB of HBM reads and writes; 35.1 ms runtime.
- **FlashAttention:** 4.4 GB of HBM reads and writes; 11.7 ms runtime.

The measurements show the systems consequence of the I/O bound: substantially less HBM traffic coincides with lower runtime even though FlashAttention performs additional recomputation. They do not imply that every configuration becomes compute bound; tile size, sequence length, and kernel implementation determine the realized regime.

C.6 Kaplan & Chinchilla Compute-Optimal Scaling Law Derivation

Determining the compute-optimal allocation between model parameters N and training dataset tokens D for a given floating-point budget C is a central decision in large language model system design. This section provides the rigorous mathematical derivation of compute-optimal scaling laws using Lagrange multipliers, contrasting the original findings of Kaplan et al. (2020) with the compute-optimal corrections of Hoffmann et al. (2022).

C.6.1 Parametric loss surface formulation

The cross-entropy evaluation loss $\mathcal{L}(N, D)$ of an autoregressive transformer language model trained with N non-embedding parameters on D tokens is modeled by the power-law formulation:

$$\mathcal{L}(N, D) = E + \frac{A}{N^\alpha} + \frac{B}{D^\beta}$$

where:

- E : Irreducible loss representing the entropy of natural language data.
- $\frac{A}{N^\alpha}$: Capacity bottleneck loss resulting from finite model parameters N .
- $\frac{B}{D^\beta}$: Dataset size bottleneck loss resulting from finite training tokens D .
- $A, B > 0$: Empirical scaling constants.
- $\alpha, \beta > 0$: Power-law scaling exponents for model size and dataset size, respectively.

C.6.2 Constrained optimization via Lagrange multipliers

The total floating-point operations required to train a transformer with N non-embedding parameters for D tokens is given by:

$$C = 6ND \text{ FLOPs}$$

where $2ND$ accounts for the forward pass and $4ND$ accounts for the backward pass.

Given a fixed training compute budget C_0 , we formulate the constrained optimization problem to minimize the reducible loss $\hat{\mathcal{L}}(N, D) = \frac{A}{N^\alpha} + \frac{B}{D^\beta}$:

$$\min_{N, D} \left(\frac{A}{N^\alpha} + \frac{B}{D^\beta} \right) \quad \text{subject to} \quad g(N, D) = 6ND - C_0 = 0$$

We construct the Lagrangian function $\mathcal{L}_{\text{Lagrange}}(N, D, \lambda)$:

$$\mathcal{L}_{\text{Lagrange}}(N, D, \lambda) = \frac{A}{N^\alpha} + \frac{B}{D^\beta} + \lambda(6ND - C_0)$$

Taking partial derivatives with respect to N , D , and λ and setting them to zero:

$$\frac{\partial \mathcal{L}_{\text{Lagrange}}}{\partial N} = -\alpha AN^{-\alpha-1} + 6\lambda D = 0 \implies \alpha AN^{-\alpha-1} = 6\lambda D \quad (1)$$

$$\frac{\partial \mathcal{L}_{\text{Lagrange}}}{\partial D} = -\beta BD^{-\beta-1} + 6\lambda N = 0 \implies \beta BD^{-\beta-1} = 6\lambda N \quad (2)$$

$$\frac{\partial \mathcal{L}_{\text{Lagrange}}}{\partial \lambda} = 6ND - C_0 = 0 \quad (3)$$

Multiplying Equation (1) by N and Equation (2) by D :

$$\alpha AN^{-\alpha} = 6\lambda ND$$

$$\beta BD^{-\beta} = 6\lambda ND$$

Equating these expressions yields the fundamental power-law optimality condition:

$$\alpha AN^{-\alpha} = \beta BD^{-\beta}$$

This condition proves that at the compute-optimal allocation, the marginal loss reduction per fractional increase in parameter count must equal the marginal loss reduction per fractional increase in dataset tokens.

C.6.3 Analytical derivation of compute-optimal exponents

From $\alpha AN^{-\alpha} = \beta BD^{-\beta}$, we solve for D in terms of N :

$$D^\beta = \left(\frac{\beta B}{\alpha A}\right) N^\alpha \implies D = \left(\frac{\beta B}{\alpha A}\right)^{1/\beta} N^{\alpha/\beta}$$

Substituting D into the compute constraint $C = 6ND$:

$$C = 6N \left(\frac{\beta B}{\alpha A}\right)^{1/\beta} N^{\alpha/\beta} = 6 \left(\frac{\beta B}{\alpha A}\right)^{1/\beta} N^{\frac{\alpha+\beta}{\beta}}$$

Solving explicitly for the compute-optimal parameter count $N^*(C)$:

$$N^*(C) = \left(\frac{\alpha A}{\beta B}\right)^{\frac{1}{\alpha+\beta}} \left(\frac{C}{6}\right)^{\frac{\beta}{\alpha+\beta}} \propto C^{\frac{\beta}{\alpha+\beta}}$$

By symmetry, solving for the compute-optimal token count $D^*(C)$:

$$D^*(C) = \left(\frac{\beta B}{\alpha A}\right)^{\frac{1}{\alpha+\beta}} \left(\frac{C}{6}\right)^{\frac{\alpha}{\alpha+\beta}} \propto C^{\frac{\alpha}{\alpha+\beta}}$$

Empirical exponents and the $D/N \approx 20$ rule

Fitting the parametric loss surface to over 400 experimental runs, Hoffmann et al. (2022) estimated $\alpha \approx 0.34$ and $\beta \approx 0.28$. Substituting these values gives exponents close to the paper's parametric-model result:

$$\frac{\beta}{\alpha+\beta} \approx 0.45, \quad \frac{\alpha}{\alpha+\beta} \approx 0.55$$

Thus the parametric model predicts:

$$N^*(C) \propto C^{0.45}, \quad D^*(C) \propto C^{0.55}$$

The paper's two direct empirical approaches report near-balanced exponents of $(0.50, 0.50)$ and $(0.49, 0.51)$, which motivate the common square-root shorthand. Under the parametric fit, however, D^*/N^* grows slowly with compute rather than remaining strictly constant. Across the regime studied, roughly 20 training tokens per parameter is a useful Chinchilla heuristic, not a universal invariant.

🔍 Systems Perspective C.5: Kaplan vs. Chinchilla scaling

The difference between Kaplan et al. (2020) and Hoffmann et al. (2022) highlights how hyperparameter choices in empirical benchmarking can alter system design conclusions:

- **Kaplan et al. (2020)** estimated $\alpha \approx 0.057$ and $\beta \approx 0.28$, predicting $N \propto C^{0.73}$ and $D \propto C^{0.27}$. This led the field to scale model sizes aggressively while holding training dataset sizes relatively small (e.g., GPT-3 175B trained on 300B tokens, giving $D/N = 1.71$).
- **Methodological difference:** Kaplan et al. (2020) used a fixed token budget and learning-rate schedule across models. Hoffmann et al. (2022) matched the schedule to each training horizon and found that the earlier setup understated the value of training smaller models on more data.
- **Hoffmann et al. (2022)** found near-balanced parameter and token scaling: its three approaches report exponent pairs of $(0.50, 0.50)$, $(0.49, 0.51)$, and $(0.46, 0.54)$ rather than one exact universal ratio.
- **Practical Impact:** Chinchilla (70B parameters trained on 1.4T tokens, $D/N = 20$) outperformed GPT-3 (175B parameters trained on 300B tokens) across downstream benchmarks while requiring $2.5\times$ less compute during inference and $2.5\times$ lower weight memory at the same precision (Hoffmann et al. 2022).

Summary

An algorithm becomes a systems workload through its shapes, operation graph, data dependencies, and numerical representation. Linear-algebra primitives, memory layouts, broadcasting, and sparse formats determine work, traffic, locality, and usable parallelism. Automatic differentiation and backpropagation add saved activations and backward computation; computational graphs then expose opportunities for fusion and recomputation. The Baur–Strassen result bounds reverse-mode overhead, while FlashAttention shows how tiling and an online recurrence can reduce HBM traffic without changing attention semantics.

The same accounting connects kernels to scaling decisions. Compute-optimal scaling turns a fixed budget into a model-data allocation, but the fitted exponents remain empirical rather than universal. Across these models, the reusable method is to expose operations, bytes, saved state, dependencies, and asymptotic constraints before selecting an optimization.

Machine Foundations

A datasheet reports peak throughput; a workload almost never reaches it. Closing that gap requires knowing which physical limit binds execution first: latency, memory bandwidth, arithmetic intensity, or energy. This appendix collects the durable hardware models and order-of-magnitude values that anchor single-node reasoning throughout the book, providing the mathematical baselines where chapter estimates can be audited or extended. The models presume an undergraduate background in computer architecture, particularly the memory hierarchy.

How to Use This Appendix

This appendix is designed as a reference. When diagnosing performance issues, use this appendix to translate a vague symptom (“it’s slow”) into a specific constraint (“memory bound at batch size one”) and then choose the lever that can actually move.

Conventions used here follow the book-wide notation (for example, B is reserved for batch size and BW for bandwidth).

- Sanity-check feasibility: Start with section D.1 for order-of-magnitude numbers.
- Diagnose the dominant ceiling: Use the Roofline Model in section D.2.1 to decide whether the workload is compute bound or memory bound.
- Reason about scaling limits: Use Amdahl’s and Gustafson’s Laws in section D.2.3 to understand why adding accelerators may not reduce time-to-train.
- Choose the right precision: Use section D.4.1 to reason about FP32 vs. BF16/FP16 vs. INT8 as a systems trade-off.
- Follow the chapter treatment: For the full narrative, use chapter 11, chapter 8, and chapter 13.

Although modern hardware encompasses diverse accelerator architectures, interconnect topologies, and packaging options, the physical laws and bottleneck models detailed here govern them all.

D.1 Numbers to Know

Just as Jeff Dean’s “Latency Numbers Every Programmer Should Know”¹ shaped a generation of systems engineers, these reference numbers provide the order-of-magnitude intuition essential for ML systems design. Although absolute values and ratios vary by technology and workload, the hierarchy is more durable. Memorize the relationships; use the specific numbers as sanity checks.

Systems Perspective D.1: Three numbers that matter most

- Energy ratio: In the cited 45 nm reference, a 32-bit DRAM read uses $\sim 581\times$ the energy of one FP16 multiply. This motivates arithmetic intensity.
- Training-state footprint: Model weights (2 bytes FP16) + gradients (2 bytes FP16) + master weights (4 bytes FP32) + optimizer states for Adaptive Moment Estimation (Adam) at 8 bytes. That totals 16 bytes per parameter, so a 7B model needs 112 GB just to start training.
- Fiber propagation limit: Light travels about 200 km/ms in fiber. The table’s ~ 40 ms cross-country round trip includes reducible overhead, but propagation cannot be optimized away.

¹ | **Jeff Dean:** A Google Senior Fellow and one of the architects of Google’s distributed systems infrastructure, including MapReduce, BigTable, and TensorFlow. His latency numbers, originally presented with Peter Norvig around 2010, became a canonical reference for systems engineers. The numbers have been updated over the years as hardware evolved, but the *hierarchy* of latencies remains remarkably stable; Colin Scott’s interactive visualization shows the latency hierarchy across hardware generations (Scott 2012).

Reference relationships

These relationships mix physical or arithmetic bounds with technology-specific measurements; only the bounds are invariant.

Speed of light tax

Table D.1 shows approximate round-trip baselines. Their propagation component is irreducible.

Table D.1: Speed of Light Reference: Light in fiber travels ~200 km/ms. These approximate round trips include reducible overhead; only propagation is irreducible.

Distance	Round-Trip Latency	Implication
Same data center	~1 ms	Distributed training feasible
Cross-country (US)	~40 ms	Reduces the remaining request budget
Cross-Atlantic	~60 ms	Further reduces the request budget
Cross-Pacific	~100 ms	Data locality is critical

Energy hierarchy

Table D.2 quantifies the energy cost of data movement vs. computation—the fundamental reason why arithmetic intensity dominates ML performance optimization.²

Table D.2: Energy Reference Ratios: The first three ratios use Horowitz’s 45 nm measurements; the L1/register ratio is an illustrative on-chip hierarchy anchor.

Relationship	Reference Ratio	Source and Interpretation
DRAM access vs. FP16 multiply	~581×	Horowitz 45 nm; wire capacitance scales with distance
FP32 vs. INT8 energy	~18×	Horowitz 45 nm; bit width determines switching energy
FP32 vs. FP16 energy	~3.4×	Horowitz 45 nm; narrower arithmetic reduces switching and datapath energy
L1 SRAM vs. register	~5×	Illustrative on-chip hierarchy anchor; distance to the arithmetic unit

Memory hierarchy

Table D.3 shows why the memory hierarchy is uneven: nearby on-package hops can differ by only a few times, while off-chip, storage, and network tiers introduce orders-of-magnitude jumps.

Table D.3: The Latency Hierarchy: Selected latency relationships in the memory hierarchy. The largest gaps come from crossing physical boundaries: off-chip memory, peripheral links, storage, and network fabric.

Relationship	Ratio	Why It Persists
Accelerator high-bandwidth memory vs. register	~1000× slower	On-chip vs. off-chip
SSD vs. register	~300,000× slower	Electrical vs. mechanical/flash
Network vs. local memory	~16× slower	Speed of light + switching
Accelerator memory BW vs. CPU↔Accelerator link	~52× faster	Architectural investment priority

Scaling laws

Table D.4 collects the arithmetic relationships that govern memory and compute requirements for training and inference.³

² | **Energy hierarchy sources:** The DRAM and arithmetic values come from Horowitz (2014) (45 nm). The L1 SRAM/register ratio is an illustrative hierarchy anchor, not a Horowitz measurement. The headline ratio compares a 32-bit DRAM access with a 16-bit floating-point multiply. Both the absolute values and the ratio vary with process, circuit, memory technology, and operand width; wire distance helps explain why movement remains expensive.

³ | **Training memory (Adam):** The 16 bytes/parameter rule assumes mixed-precision training with Adam. ZeRO optimization can reduce per-accelerator memory by sharding optimizer states across accelerators, but the total memory across all accelerators remains ~16× parameters.

Table D.4: Scaling Rules: These are arithmetic, not hardware-specific. Training memory includes FP16 weights (2 bytes), FP16 gradients (2 bytes), FP32 master weights (4 bytes), and Adam optimizer states (8 bytes for momentum + variance).

Rule	Formula	Example
Inference memory (FP16)	$2 \text{ bytes} \times \text{parameters}$	7B params $\rightarrow$ 14 GB
Inference memory (INT8)	$1 \text{ byte} \times \text{parameters}$	7B params $\rightarrow$ 7 GB
Training memory (Adam)	$16 \text{ bytes} \times \text{parameters}$	7B params $\rightarrow$ 112 GB
Inference FLOPs (transformer)	$\sim 2 \times \text{parameters per token}$	7B model $\rightarrow \sim 14$ GFLOPs/token
Training FLOPs	$\sim 6 \times \text{parameters} \times \text{tokens}$	7B on 1T tokens $\rightarrow 4 \times 10^{22}$ FLOPs

Illustrative latency budgets

Latency budgets depend on application-specific safety requirements, user expectations, and service-level objectives. Table D.5 lists illustrative targets rather than universal limits.

Table D.5: Illustrative Latency Targets: Production requirements depend on the application and its service-level objectives.

Application	Budget	Constraint
Autonomous braking	<10 ms	At 100 km/h, 10 ms = 28 cm of travel
Voice assistant	<100 ms	Human perception of “instant”
Web search	<200 ms	User patience threshold
Video streaming	<1 s	Buffer tolerance
Batch training	hours–days	Throughput dominates latency

Current hardware reference (c. 2024)

These numbers reflect the current generation. Use them for back-of-envelope calculations and recheck them against current hardware.

Memory latency and bandwidth

Table D.6 provides representative latency and bandwidth anchors across the memory hierarchy.

Table D.6: Memory Hierarchy (c. 2024): Representative latency and bandwidth anchors across registers, caches, memory, interconnects, networks, and storage. Together they show why locality and data movement shape performance.

Level	Latency	Bandwidth
Register	$\sim 0.3 \text{ ns}$	—
L1 Cache	$\sim 1 \text{ ns}$	—
L2 Cache	$\sim 4 \text{ ns}$	—
GPU HBM3	$\sim 300 \text{ ns}$	3.4 TB/s
PCIe Gen5 (CPU $\leftrightarrow$ GPU)	$\sim 1000 \text{ ns}$	64 GB/s
CPU DRAM	$\sim 100 \text{ ns}$	50 GB/s
InfiniBand (network)	$\sim 5000 \text{ ns}$	50 GB/s
NVMe SSD	$\sim 100000 \text{ ns}$	7 GB/s

Compute throughput

Table D.7 compares representative peak throughput and power for data-center GPUs and a mobile NPU.

Table D.7: Compute Reference (c. 2024): Representative peak throughput and power for data-center GPUs and a mobile NPU. Because the rows use different precision modes and operation conventions, compare them only within a matched workload and precision context.

Platform	FP16/BF16	INT8/FP8-class	Power
Data center GPU (H100)	989 TFLOP/s	1979 TFLOP/s (FP8 peak)	700 W
Data center GPU (A100)	312 TFLOP/s	624 TOPS	400 W
Mobile NPU	—	35 TOPS	3–5 W

Roofline ridge points

Table D.8 defines the arithmetic intensity thresholds that determine whether a workload is memory bound or compute bound.

Table D.8: Arithmetic Intensity Thresholds (c. 2024): Batch-1 weight-streaming inference is often memory bound; larger batches may cross the ridge.

Accelerator	Ridge Point	Implication
A100 (FP16)	153 FLOP/byte	Below → memory-bound
H100 (FP16)	295 FLOP/byte	Higher bar for compute-bound

Systems Perspective D.2: A note on terminology: GPUs and accelerators

Throughout this book, “accelerator” is the general term for hardware acceleration. The principles—roofline analysis, memory hierarchies, numerical precision, and performance modeling—apply to GPUs, Tensor Processing Units (TPUs), NPUs, application-specific integrated circuits, and other specialized AI accelerators. Vendor-specific features such as CUDA and NVLink are identified explicitly.

Knowing the numbers is only the first step. The real power comes from compact models that reveal *which* number matters for the bottleneck at hand. The Roofline Model begins that diagnosis by translating raw hardware specs into actionable performance ceilings.

D.2 Physics of Computing

Raw hardware specs—TFLOP/s, TB/s, watt budgets—are necessary but insufficient for performance reasoning. Without compact analytical models, an engineer cannot distinguish a compute-bound workload from a memory-bound one, or predict whether doubling GPUs will halve training time. The models in this section turn those distinctions into quantitative checks.

D.2.1 The Roofline model

The Roofline Model (Williams et al. 2009) bounds how fast a workload can run on a given hardware target. The answer depends on whether the workload runs out of compute or memory bandwidth first.

Every operation has an **arithmetic intensity**: the ratio of computations performed to bytes moved from memory. Matrix multiplication has high arithmetic intensity because each loaded element is reused many times. Element-wise operations like rectified linear unit (ReLU) have low intensity because each operation loads a number, performs one computation, and writes it back. As figure D.1 illustrates, each workload is bounded by either memory bandwidth or compute throughput, and its arithmetic intensity determines which ceiling it hits first.




Figure D.1: The Roofline Model: The sloped memory-bandwidth ceiling and horizontal peak-compute ceiling meet at the ridge point. Plotting a workload by arithmetic intensity identifies which roofline ceiling applies.

The **ridge point** determines the hardware’s balance. If a workload’s intensity falls below this point, it is memory-bound (sloped region). If above, it is compute-bound (flat region).

$$\text{Arithmetic Intensity} = \frac{\text{FLOPs}}{\text{bytes accessed}} \quad \text{Ridge Point} = \frac{\text{Peak FLOP/s}}{\text{Memory Bandwidth}}$$

🔍 Systems Perspective D.3: Batch size controls arithmetic intensity

For matrix multiplications, arithmetic intensity scales with the batch dimension. For $\mathbf{Y} = \mathbf{XW}$, where $\mathbf{X}$ is $(B \times d_{\text{in}})$ and $\mathbf{W}$ is $(d_{\text{in}} \times d_{\text{out}})$:

- FLOPs: $2 \times B \times d_{\text{in}} \times d_{\text{out}}$ (multiply-adds)
- Bytes: Moving X , W , and Y costs approximately $s_{\text{elem}}(Bd_{\text{in}} + d_{\text{in}}d_{\text{out}} + Bd_{\text{out}})$ bytes.

Doubling B doubles FLOPs while weight traffic stays nearly constant when weights dominate, increasing arithmetic intensity. This is why batching often helps inference serving, although activation traffic, cache behavior, shape, precision, and the ridge point determine whether it reaches the compute ceiling.

D.2.1.1 A concrete example: The A100 analysis

Consider an NVIDIA A100 GPU with FP16 Tensor Core performance of 312 TFLOP/s and HBM2e bandwidth of 2.04 TB/s. The ridge point is 312 TFLOP/s / 2.04 TB/s = 153 FLOP/byte (the Tera prefixes cancel, yielding FLOP/byte).

Two common operations fall on opposite sides of that ridge. General matrix multiply (GEMM) on two square matrices of size 4096 by 4096 has arithmetic intensity of approximately 1365 FLOP/byte, so 1365 FLOP/byte > 153 FLOP/byte and the operation is compute bound. An element-wise ReLU on an FP16 tensor performs 1 comparison while reading 2 bytes and writing 2 bytes (4 bytes total), yielding an arithmetic intensity of only $1/4 = 0.25$ FLOP/byte. Because 0.25 FLOP/byte $\ll$ 153 FLOP/byte, the operation is severely memory bound, achieving only about 0.16 percent of peak TFLOP/s. This contrast explains why modern frameworks fuse operations. Combining ReLU directly with the preceding MatMul kernel avoids writing intermediate results to memory, effectively preserving memory bandwidth.

D.2.2 Dimensional analysis

The Roofline Model helps diagnose *where* a bottleneck lies. Before applying any performance equation, however, it must be verified as physically meaningful. Dimensional analysis provides this

sanity check: any valid equation must be **dimensionally homogeneous**—every term must resolve to the same units. If they do not, the equation contains an error.

For serialized phases, consider the iron law of ML systems (principle 3) introduced in section 1.7:

$$T = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}}$$

Correctness is verified by confirming that every term resolves to time (seconds):

$$T[s] = \underbrace{\frac{D_{\text{vol}}[\text{bytes}]}{\text{BW}[\text{bytes/s}]}}_{\text{seconds}} + \underbrace{\frac{O[\text{FLOPs}]}{R_{\text{peak}}[\text{FLOP/s}] \cdot \eta_{\text{hw}}[1]}}_{\text{seconds}} + \underbrace{L_{\text{lat}}[s]}_{\text{seconds}}$$

- Data term: $\frac{\text{bytes}}{\text{bytes/s}} = \text{bytes} \times \frac{\text{s}}{\text{bytes}} = \text{s}$
- Compute term: $\frac{\text{FLOPs}}{\text{FLOP/s}} = \text{FLOPs} \times \frac{\text{s}}{\text{FLOP}} = \text{s}$
- Latency term: Already in seconds.

The equation is physically consistent. Apply this technique to any systems equation: if the dimensions do not match, the formula is wrong. “FLOPs” and “Bandwidth” cannot be traded directly because they have different units. Any such trade-off must convert through Time, which is precisely what the iron law quantifies.

Once an equation is dimensionally sound, the next question is how its time terms scale across devices. Section D.2.3 answers that question with two complementary bounds.

D.2.3 Amdahl’s Law and Gustafson’s Law

Parallelization is the primary tool for scaling ML, but its limits depend on *how* the workload scales. These two laws frame the fundamental tension in parallel computing. **Amdahl’s law** is the pessimist’s view, governing how much faster a *fixed* task can run (optimizing latency). **Gustafson’s law** is the optimist’s view, governing how much *more* work can be done in the same time (optimizing throughput).

D.2.3.1 Strong scaling (Amdahl’s Law)

Strong scaling measures how much faster a fixed-size problem runs as processors are added.

Amdahl’s Law (Amdahl 1967) states that the speedup is limited by the serial portion of the task.⁴ If a fraction s of the task is serial (cannot be parallelized) and $p = 1 - s$ is parallelizable, the maximum speedup with n processors is:

$$\text{Speedup}(n) = \frac{1}{s + \frac{1-s}{n}}$$

As $n \rightarrow \infty$, the term $\frac{1-s}{n} \rightarrow 0$, and the speedup converges to $1/s$.

To see Amdahl’s Law in action, suppose 5 percent of a training step is serial overhead (for example, Python global interpreter lock (GIL), kernel launch latency) and 95 percent is parallelizable matrix math:

- With $n = 1$, speedup is 1.
- With $n = 8$, speedup is $1/(0.05 + 0.95/8) \approx 5.9\times$.
- With $n \rightarrow \infty$, speedup is capped at $1/0.05 = 20\times$.

No matter how many accelerators are added, this fixed workload cannot run faster than $20\times$.

D.2.3.2 Weak scaling (Gustafson’s Law)

Weak scaling measures how much larger a problem can become while holding runtime fixed as processors are added.⁵

This is the reality of Large Language Models. Rather than using 1,000 accelerators to train a model on a small dataset in milliseconds, they are used to train on a dataset $1,000\times$ larger in reasonable time.

⁴ | **Gene Amdahl (1922–2015):** A legendary computer architect at IBM, where he was the chief architect of the System/360. He later founded Amdahl Corporation to compete with IBM in the mainframe market.

⁵ | **John Gustafson:** A computer scientist known for his work in parallel computing and for introducing the Unum (universal number) format. His law was a direct response to the perceived “limits” of Amdahl’s Law when applied to massive scale.

For n processors, let s be the serial fraction of execution time measured on those processors for the scaled workload. Gustafson's Law (Gustafson 1988) models this "scaled speedup":

$$\text{Scaled Speedup}(n) = n - s(n - 1)$$

Because the parallel part grows linearly with n while s remains fixed, weak scaling can keep efficiency high as the problem grows.

Using the same 5 percent serial overhead ($s = 0.05$), Gustafson's Law tells a very different story:

- With $n = 1$, speedup is 1.
- With $n = 8$, Scaled Speedup is $8 - 0.05 \times (7) = 8 - 0.35 = 7.65\times$.
- With $n = 1000$, Scaled Speedup is $1000 - 0.05 \times (999) \approx 950\times$.

In weak scaling, efficiency remains high because the useful work (training the model) scales up to dwarf the fixed overheads.

The same scaling lens turns model size, data size, hardware count, and utilization into a single training-time estimate.

Napkin Math D.1: The training time equation

Problem: How long does dense-transformer training take for a stated model size, token count, accelerator count, peak rate, and utilization?

Math:

$$T \approx \frac{6 \cdot P \cdot D}{N_{\text{GPU}} \cdot R_{\text{peak}} \cdot \eta_{\text{hw}}}$$

Variables:

- Training FLOP factor: The factor 6 derives from $2P$ FLOP/token for the forward pass and $4P$ FLOP/token for the backward pass ($2P$ for activation gradients and $2P$ for weight gradients), yielding $6PD$ FLOP total across D training tokens.
- P : Number of model parameters.
- D : Number of training tokens.
- N_{GPU} : Number of GPUs.
- R_{peak} : Peak FLOP/s of one accelerator.
- η_{hw} : Hardware utilization, typically 30 percent–50 percent in this training estimate.

Example: Training a 1B parameter model on 20B tokens using 1 A100 (312 TFLOP/s) at 40 percent utilization. Total FLOPs = $6 \times 1 \times 10^9 \times 2 \times 10^{10} = 1.2 \times 10^{20}$ FLOPs. Throughput = $1 \times (3.12 \times 10^{14}) \times 0.40 \approx 1.248 \times 10^{14}$ FLOP/s. $T = \frac{1.2 \times 10^{20}}{1.248 \times 10^{14}} \approx 961,538.5$ seconds $\approx 16,025.6$ minutes.

Systems insight: Under these assumptions, training takes 961,538.5 seconds (≈ 16025.6 minutes, or about 11.1 days); the estimate makes model scale, data scale, effective accelerator throughput, and utilization explicit levers.

D.2.4 Little's Law

For capacity planning in stable inference systems, **Little's law** (Little 1961) relates long-run mean concurrency (Q_{req}), mean arrival rate (λ_{arr}), and mean time in system (T_{lat}):⁶

$$Q_{\text{req}} = \lambda_{\text{arr}} \times T_{\text{lat}}$$

To see this in practice, consider sustaining 1,000 QPS with 50 ms average latency. The law dictates that the system must support $1000 \times 0.05 \text{ s} = 50$ concurrent requests.

This sizes worker pools, but memory needs another assumption. If every in-system request simultaneously holds 1 GB of device state, the average population implies 50 GB; queued requests may instead stay in host memory or share batched state. If 24 device slots each remain occupied for

⁶ | **John Little:** An Institute Professor at MIT and a pioneer in the field of operations research. His law, proved in 1961, is fundamental to queuing theory and is used across fields from manufacturing to computer network analysis.

50 ms, their throughput proxy is $N_{\text{max}}/T_{\text{lat}} = 24/0.05 = 480$ QPS; Little’s Law alone sets no physical maximum.

These physics-based models—Roofline, Amdahl, Gustafson, and Little—diagnose *where* bottlenecks lie. Translating those diagnoses into actionable optimizations, however, requires understanding the concrete hardware structures that impose them: caches, memory buses, and interconnects.

D.3 Computer Architecture Essentials

A GPU advertises 1,000 TFLOP/s, yet a kernel achieves only 30 TFLOP/s. The gap may reflect data movement, insufficient parallelism, synchronization, launch overhead, or software inefficiency. While physics sets theoretical performance bounds, computer architecture defines the machinery that determines how close a real workload can get. The following discussion covers the latency, bandwidth, and energy trade-offs that shape system design.

D.3.1 Latencies every programmer should know

The first step in systems intuition is understanding the cost of distance. Table D.9 quantifies how long the processor waits for data from different levels of the memory hierarchy. If accessing a register is like picking up a nearby pencil, fetching from high-bandwidth memory (HBM) is walking across the office, and fetching from disk is flying to the moon.

Table D.9: The Latency Hierarchy: Representative access times for modern AI hardware; the approximate cycle counts use a 3 GHz reference clock. The jump from SRAM cache to HBM makes cache locality a major performance factor.

Component	Latency (ns)	Cycles (Approx)	Relative “Distance”
Register	~0.3 ns	1 cycle	10 seconds
L1 Cache	~1 ns	3–4 cycles	33.3 seconds
L2 Cache	~4 ns	12 cycles	2.2 minutes
HBM3 (GPU Memory)	~300 ns	1,000 cycles	2.8 hours
NVLink (GPU-GPU)	~500 ns	1,500 cycles	4.6 hours
PCIe (CPU-GPU)	~1000 ns	3,000 cycles	9.3 hours
InfiniBand (Network)	~5000 ns	15,000 cycles	1.9 days
SSD (NVMe)	~100000 ns	300,000 cycles	38.6 days

D.3.2 The AI hardware cheat sheet (modern reference)

Latency measures the wait for the *first* byte; bandwidth measures how many bytes follow. Table D.10 provides the constants for back-of-the-envelope “Roofline” calculations. These represent the “standard units of compute” for the current era of machine learning.

Table D.10: Reference Specs: Key constants for quantitative analysis. Always check specific datasheets; HBM capacity follows vendor units, and interconnect bandwidth is aggregate bidirectional, so topology and direction matter.

Spec	NVIDIA H100 (SXM)	Google TPU v5p	System Impact
BF16 Peak	989 TFLOP/s	459 TFLOP/s	The “Speed Limit” (R_{peak})
Memory Bandwidth	3.35 TB/s	2.76 TB/s	The “Width of the Pipe” (BW)
HBM Capacity	80 GB	95 GiB	Max Model Size (P)/Batch Size (B)
On-chip SRAM	50 MB L2	128 MiB VMEM/chip	Critical for Operator Fusion
Interconnect	900 GB/s (NVLink)	1200 GB/s (Inter-Chip Interconnect, ICI)	Determines Model Parallelism Scaling

D.3.3 The memory hierarchy

Computer systems use a hierarchy because no single technology provides both high capacity and low latency. Figure D.2 shows this trade-off: keeping data higher in the pyramid (registers/cache) reduces access latency when data movement is on the critical path.




Figure D.2: The Memory Hierarchy: Performance depends on data proximity. Accessing HBM is roughly $1,000\times$ slower than registers; accessing SSD is roughly $300,000\times$ slower.

The memory hierarchy is the fundamental physical constraint of machine learning systems. Table D.11 consolidates the physical properties—latency, bandwidth, and energy—across the entire stack.

Table D.11: Physical Properties of the Memory Hierarchy (c. 2024): Representative latency and bandwidth anchors across the memory hierarchy. The hierarchy spans roughly five orders of magnitude in latency. NVLink is aggregate bidirectional; only the cited 45 nm DRAM energy is shown because other tiers are implementation-specific. Keeping data higher in the hierarchy can therefore deliver large performance gains.

Layer	Technology	Latency	Bandwidth	Energy Reference (per 32b)
Registers	Flip-Flops	~ 0.3 ns	—	—
L1 Cache	SRAM	~ 1 ns	—	—
L2 Cache	SRAM	~ 4 ns	—	—
Memory (Local)	HBM3	~ 300 ns	3350 GB/s	640 pJ
Interconnect	NVLink 4.0	~ 500 ns	900 GB/s	—
Host Link	PCIe Gen5	~ 1000 ns	64 GB/s	—
System RAM	DDR5	~ 100 ns	50 GB/s	—
Network (Fabric)	InfiniBand NDR	~ 5000 ns	50 GB/s	—
Storage (Local)	NVMe SSD	~ 100000 ns	7 GB/s	—

The hierarchy’s energy costs reveal why data movement dominates modern system design.

Systems Perspective D.4: The high cost of data movement

In the cited 45 nm reference, fetching a 32-bit value from DRAM costs roughly $581\times$ the energy of one FP16 multiply (~ 640 pJ vs. ~ 1.1 pJ). The ratio is platform-specific and does not alone determine workload energy, which also depends on access counts and reuse. It nevertheless explains why arithmetic intensity matters.

D.3.4 Bandwidth vs. latency

Bandwidth (throughput) and latency (delay) are distinct constraints. For data volume D_{vol} sent over effective bandwidth BW with fixed path latency L_{lat} , and assuming the latency and serialization costs do not overlap, total transfer time is:

$$T = L_{\text{lat}} + \frac{D_{\text{vol}}}{BW}$$

The crossover point is the transfer size where the two terms are equal:

$$D_{\text{vol,cross}} = L_{\text{lat}} \times \text{BW}$$

For transfers below $D_{\text{vol,cross}}$ (for example, small request or control messages), latency dominates. For transfers above it (for example, loading weights), bandwidth dominates.

Consider sending data over a 10 Gb/s link with 10 ms base request/response latency. The crossover size is 12.5 MB, so the dominant term depends on which side of that threshold the transfer falls.

- Latency-bound packet (1 KB):
 - Transmission: $1 \text{ KB} \times 8/10 \text{ Gb/s} \approx 0.8 \text{ } \mu\text{s}$.
 - Total Time $\approx 10 \text{ ms} + 0.8 \text{ } \mu\text{s} \approx 10 \text{ ms}$.
 - Result: Fixed path and protocol latency dominates; propagation is one component.
- Bandwidth-bound checkpoint (1 GB):
 - Transmission: $1 \text{ GB} \times 8/10 \text{ Gbps} \approx 800 \text{ ms}$.
 - Total Time $\approx 10 \text{ ms} + 800 \text{ ms} = 810 \text{ ms}$.
 - Result: Base latency is negligible; bandwidth is the bottleneck.

Architecture determines how fast data can move, but numerical precision directly controls *how much* data must move. Halving precision from FP32 to FP16 halves bytes per parameter and can nearly double useful bandwidth when the hardware, kernels, and model accuracy support it. Understanding these trade-offs requires a closer look at how numbers are represented in hardware.

D.4 Numerical Representations

Statistics characterizes data distributions; numerical representations determine how the system stores those values. In ML systems, the choice of precision (FP32 vs. BF16 vs. INT8) is a direct trade-off between statistical fidelity and hardware throughput.

D.4.1 Floating-point format comparison

IEEE 754 formats such as FP32 and FP16, together with AI-specific formats such as BF16 and FP8, define different trade-offs between dynamic range (the span of representable values) and precision (the granularity of values within that range). Table D.12 summarizes the key formats and their use cases, while figure D.3 visualizes the bit allocations.

Table D.12: Numerical Format Comparison: Each format trades off precision, dynamic range, memory footprint, and compute throughput. BF16 shares FP32's normal exponent range while using half the storage.

Format	Bits	Exponent	Mantissa	Dynamic Range	Typical Use Case
FP32	32	8	23	$\sim 10^{-38}$ to 10^{38}	Training (full precision), reference inference
FP16	16	5	10	Normal: 6.10×10^{-5} – 6.55×10^4	Training often with loss scaling, inference
BF16	16	8	7	FP32 normal exponent range	Training, usually without loss scaling
FP8	8	4 or 5	3 or 2	Varies	Training/inference on supported hardware
INT8	8	N/A	N/A	-128 to 127	Inference after quantization

Beyond bit width, the allocation of bits between exponent and mantissa determines what range of values each format can represent.

Among these formats, BF16⁷ deserves special attention (Wang and Kanwar 2019; Kalamkar et al. 2019). It occupies the same 16-bit storage class as FP16 but behaves differently in training, so nominal width alone does not predict the operational burden. The following systems perspective connects bit allocation to memory traffic, numerical stability, and deployment efficiency.

⁷ **BF16 (brain floating point 16):** Originally introduced with Google TPuv2 and later adopted by Intel, Arm, and NVIDIA, BF16 became a common training format across accelerator ecosystems. Its adoption made format support in hardware, kernels, and distributed-training libraries a portability concern rather than a vendor-specific feature.




Figure D.3: Numerical Format Bit Layouts: A visual comparison of bit allocations. BF16 preserves FP32’s 8-bit exponent and therefore its normal exponent range. FP16 trades range for precision and often uses loss scaling to reduce underflow.

🔍 Systems Perspective D.5: The dynamic range wall

The choice of numerical format is a direct application of the iron law of ML systems (principle 3). Reducing precision from FP32 to BF16 or FP16 halves the data volume term, potentially doubling throughput on memory-bound workloads. However, the *type* of 16-bit format determines the engineering complexity:

- **Dynamic range (the exponent):** BF16 preserves the eight-bit exponent of FP32 and therefore the same normal exponent range (Kalamkar et al. 2019).
- **Precision (the mantissa):** FP16 has a larger 10-bit mantissa than BF16 (7 bits), offering higher precision for values within its range. Its five-bit exponent, however, raises underflow risk. FP16 training therefore often uses **loss scaling**, an operational step that multiplies gradients by a large constant to keep more values representable (Micikevicius et al. 2017).
- **Energy efficiency:** INT8 operations can be substantially more efficient than floating-point equivalents because they move less data and use simpler integer arithmetic paths. Moving to INT8 for inference is a primary lever for deploying neural networks under tight memory, latency, or power budgets (Jacob et al. 2018; Krishnamoorthi 2018).

D.4.2 Integer quantization

Quantization maps continuous floating-point values to discrete integers, typically INT8. The key challenge is choosing how to map the floating-point range to integers. Two approaches dominate.

For symmetric quantization, let $s_{\text{quant}} = \max |x|$ over the calibrated range and assume x is clipped to $[-s_{\text{quant}}, s_{\text{quant}}]$. The signed INT8 code is:

$$x_{\text{int}} = \text{round} \left(\frac{x}{s_{\text{quant}}} \times 127 \right)$$

This mapping works well for weight distributions centered around zero.

Asymmetric quantization handles distributions that are not centered (common after ReLU, which produces only nonnegative values) by shifting the range before scaling. Let $x_{\min}$ and $x_{\max}$ bound the calibrated range, set the normalization span $s_{\text{quant}} = x_{\max} - x_{\min}$, and clip x to $[x_{\min}, x_{\max}]$. A common unsigned 8-bit mapping is:

$$x_{\text{uint8}} = \text{round} \left(\frac{x - x_{\min}}{s_{\text{quant}}} \times 255 \right)$$

The choice between symmetric and asymmetric quantization depends on the tensor’s distribution and has measurable accuracy implications.

Summary

Machine-side diagnosis begins with units, ceilings, and working sets. Reference numbers establish orders of magnitude; the Roofline Model and the iron law identify whether operations, bytes, or latency bind; Amdahl's Law, Gustafson's Law, and Little's Law test speedup, scaling, and concurrency claims; and dimensional analysis catches impossible results. The memory hierarchy and interconnect show how those limits appear as latency and bandwidth constraints, while numerical representations show how dynamic range, precision, and quantization change storage, energy, throughput, and accuracy.

Treat peak specifications as ceilings, not predictions. Compare a measurement with the relevant ceiling, then trace the remaining gap to utilization, data movement, runtime overhead, architecture, or software. A back-of-envelope estimate is useful when it states its assumptions and units and makes the next measurement obvious.

System Assumptions

A napkin calculation is only as trustworthy as the inputs behind it, and those inputs are easy to lose track of: an accelerator’s peak rate, a model’s parameter count, an electricity price, an energy cost per operation. This appendix gathers those canonical values and unit conventions into a single reference sheet, allowing any quantitative estimate in the book to be audited, updated with newer figures, or checked across chapters. It also fixes the unit conventions those estimates depend on. The chapters can then argue from numbers without restating where each one came from.

How to Use This Appendix

When a chapter says an H100 delivers a certain ridge point, or that training a 7B model needs a certain amount of memory, the underlying bandwidths, capacities, and model sizes are listed here. A local estimate can use the relevant Value and Unit entries. The tables are grouped by topic: accelerators, reference models, energy per access, interconnect bandwidth, economics, production-scale anchors, and unit conventions. Assumptions come from vendor datasheet peaks, published studies or industry reports, illustrative market or grid statistics, and book conventions; section E.1 lists the provenance in one place.

The calculations below expose which assumptions dominate each estimate, so a changed hardware specification or price can propagate without changing the method.

The constants in this appendix support quick, reproducible calculations. The three examples below apply them to performance, capacity, and cost.

Napkin Math E.1: Classify a roofline regime

Problem: For the stated H100 and workload intensities, which operations are compute bound and which are memory bound?

Variables: For an H100 at FP16/BF16 peak, use 989 TFLOP/s peak compute and 3.35 TB/s memory bandwidth.

Math: Divide peak FLOP/s by memory bandwidth to get the roofline ridge point: $989 \text{ TFLOP/s} / 3.35 \text{ TB/s} \approx 295.2 \text{ FLOP/byte}$. A large general matrix multiply (GEMM) with $n = 4,096$ has intensity $n/3 \approx 1,365.3 \text{ FLOP/byte}$.

Result: Operations above 295.2 FLOP/byte are compute bound; operations below it are memory bound. The GEMM example is compute bound, while a single-token autoregressive decode with intensity $\approx 1 \text{ FLOP/byte}$ is deeply memory bound.

Systems insight: The same accelerator can be compute rich and memory constrained. Arithmetic intensity determines which resource the workload actually consumes.

Throughput is only one constraint. Training also requires enough capacity for weights, gradients, and optimizer state.

Napkin Math E.2: Estimate training-state memory

Problem: How much model-state memory is required to train a 7B model with mixed-precision Adaptive Moment Estimation (Adam)?

Variables: Mixed-precision Adam stores BF16 weights, gradients, FP32 master weights, momentum, and variance. The model-state budget is 16 bytes per parameter.

Math: For 7B parameters, model state is $7 \times 10^9 \times 16 \text{ bytes} = 112 \text{ GB}$.

Result: An H100 has 80 GB of high-bandwidth memory (HBM), so the model state alone exceeds a single accelerator before accounting for activations.

Systems insight: Optimizer state, not weights alone, sets the floor for training memory. Capacity planning that starts from parameter bytes underestimates the real requirement.

Fitting the model in memory establishes feasibility, but not operating cost. Power and elapsed time provide the remaining quantities for a first electricity estimate.

Napkin Math E.3: Estimate accelerator electricity cost

Problem: What is the accelerator-only electricity cost of the editorial GPT-3-scale scenario?

Variables: Use 1,024 A100s, 400 W per accelerator, ~25 days wall-clock, and \$0.12/kWh electricity.

Math: The run requires roughly 3.14×10^{23} FLOPs. The A100-equivalent electricity-cost estimate is $\sim 25 \text{ days wall-clock} \times 1,024 \text{ A100s} \times 24 \text{ h/day} \times 0.4 \text{ kW} \times \$0.12/\text{kWh} = \sim \$29,491.2$.

Result: The accelerator-only electricity cost is approximately $\sim \$29,491.2$. This accelerator thermal design power (TDP)-only estimate excludes host CPUs, networking, storage, cooling, and facility overhead. The original GPT-3 run used V100-era infrastructure; this is an editorial A100-equivalent scenario, not a reported configuration or a duration derived from peak FLOP/s alone.

Systems insight: Energy is measurable from power and time, but full cost accounting must also include capital utilization, networking, storage, staffing, and failed or repeated runs.

Accelerator Specifications

These are the **accelerator assumptions** used in the book's roofline, training-memory, energy, and cost examples: peak throughput, memory bandwidth, memory capacity, and TDP for each generation the chapters cite. Values are vendor datasheet peaks—ceilings for napkin math, not sustained utilization—drawn from NVIDIA product documentation and IEEE Micro architecture articles (NVIDIA Corporation 2017, 2020b, 2020a, 2021a, 2018, 2024; Choquette et al. 2021; Choquette 2023), AMD MI300X documentation (AMD 2023), Google Tensor Processing Unit (TPU) publications (Jouppi et al. 2023), and Google Cloud's current TPU v6e specification page (Google Cloud 2026d). Capacity rows preserve official decimal-GB nameplates. Tables progress from table E.2 (Volta) and table E.3 (Turing) through current and forward-looking generations.

Master accelerator spec matrix

The master accelerator matrix in table E.1 synthesizes specifications across key GPU, TPU, Wafer-Scale, and domain-specific AI accelerators.

Table E.1: Master Accelerator Specification Matrix: Comprehensive reference matrix comparing production AI hardware across architectures, process nodes, TDP, multi-precision compute throughputs (FP64 to INT8), memory capacity and bandwidth, on-chip SRAM capacity, and FP16 Roofline Ridge Point (I_{ridge} = Peak FP16 TFLOP/s/Memory Bandwidth (TB/s)).

Architecture/Chip	Node (nm)	TDP (W)	Compute Peak TFLOPs/TOPS	Memory Capacity & Bandwidth	On-Chip SRAM	I_{ridge} (FLOP/B)
NVIDIA V100 (SXM2 32 GB)	12nm (TSMC)	300	FP64: 7.8, FP32: 15.7, FP16: 125 (TC), INT8: 125 (TC)	32 GB HBM2 @ 0.90 TB/s	16 MB	138.9
NVIDIA A100 (80 GB SXM4)	7nm N7 (TSMC)	400	FP64: 9.7/19.5, FP32: 19.5, TF32: 156, FP16/BF16: 312, INT8: 624	80 GB HBM2e @ 2.04 TB/s	40 MB	153.0
NVIDIA H100 (80 GB SXM5)	4nm 4N (TSMC)	700	FP64: 34/67, FP32: 67, TF32: 494, FP16/BF16: 989, FP8: 1,978, INT8: 1,978	80 GB HBM3 @ 3.35 TB/s	50 MB	295.2
NVIDIA H200 (141 GB SXM5)	4nm 4N (TSMC)	700	FP64: 34/67, FP32: 67, TF32: 494, FP16/BF16: 989, FP8: 1,978, INT8: 1,978	141 GB HBM3e @ 4.80 TB/s	50 MB	206.0
NVIDIA B200 (192 GB SXM)	4nm 4NP (TSMC)	1000	FP64: 45/90, FP32: 90, TF32: 1,250, FP16/BF16: 2,250, FP8: 4,500, INT8: 4,500	192 GB HBM3e @ 8.00 TB/s	128 MB	281.3
NVIDIA L40S (48 GB PCIe)	4nm 4N (TSMC)	350	FP64: 1.4, FP32: 91.6, TF32: 183, FP16/BF16: 366, FP8: 733, INT8: 733	48 GB GDDR6 @ 0.86 TB/s	96 MB	423.6
Google TPU v1	28nm (TSMC)	75	INT8: 92 TOPS (FP formats N/A)	8 GB DDR3 @ 0.034 TB/s	28 MB	N/A (2,705.9 INT8)
Google TPU v2	16nm (TSMC)	280 (board)	BF16/FP16: 45 TFLOPs	16 GB HBM @ 0.60 TB/s	32 MB	75.0
Google TPU v3	16nm (TSMC)	450 (board)	BF16/FP16: 123 TFLOPs	32 GB HBM2 @ 0.90 TB/s	32 MB	136.7
Google TPU v4	7nm N7 (TSMC)	170	BF16/FP16: 275, FP32: 275, INT8: 275	32 GB HBM2 @ 1.20 TB/s	48 MB	229.2
Google TPU v5e	7nm N7 (TSMC)	140	BF16/FP16: 197, FP8: 394, INT8: 394	16 GB HBM2e @ 0.82 TB/s	32 MB	240.5
Google TPU v5p	4nm 4N (TSMC)	450	BF16/FP16: 459, FP8: 918, INT8: 918	95 GB HBM3 @ 2.76 TB/s	96 MB	166.3
Google TPU v6e (Trillium)	4nm 4N (TSMC)	400	BF16/FP16: 926, FP8: 1,852, INT8: 1,852	32 GB HBM3 @ 1.64 TB/s	64 MB	564.6
Cerebras WSE-3	5nm N5 (TSMC)	23,000	FP16/BF16: 125,000, FP8: 250,000, INT8: 250,000	44,000 GB SRAM @ 21,000 TB/s	44,000 MB	5.95
AMD Instinct MI300X	5nm+6nm 3D	750	FP64: 163.4, FP32: 163.4, TF32: 653.7, FP16/BF16: 1,307.4, FP8: 2,614.9	192 GB HBM3 @ 5.30 TB/s	256 MB	246.7
Intel Habana Gaudi 2	7nm N7 (TSMC)	600	FP32: 38.5, TF32: 154, FP16/BF16: 307, FP8: 614, INT8: 614	96 GB HBM2e @ 2.45 TB/s	48 MB	125.3
Intel Habana Gaudi 3	5nm N5 (TSMC)	900	FP32: 115, TF32: 460, FP16/BF16: 1,835, FP8: 3,670, INT8: 3,670	128 GB HBM3e @ 3.70 TB/s	96 MB	495.9

NVIDIA V100

Table E.2: NVIDIA V100 (Volta): Peak specs from NVIDIA's architecture whitepaper and product datasheet (NVIDIA Corporation 2017, 2020b). The V100 introduced Tensor Cores for mixed-precision training and anchors baseline-generation comparisons.

Assumption	Value	Unit
Peak FP16 tensor throughput (V100)	125	TFLOP/s
Peak FP32 throughput (V100)	15.7	TFLOP/s
HBM bandwidth (V100)	900	GB/s

Assumption	Value	Unit
HBM capacity (V100)	32	GB
TDP (V100)	300	W

NVIDIA T4

Table E.3: NVIDIA T4 (Turing): Peak specs from NVIDIA Tesla T4 product documentation (NVIDIA Corporation 2018). The 70 W TDP makes this low-power inference accelerator a useful cost-per-inference anchor.

Assumption	Value	Unit
Peak FP16 tensor throughput (T4)	65	TFLOP/s
Peak INT8 throughput (T4)	130	TOPS
Memory bandwidth (T4)	320	GB/s
TDP (T4)	70	W

NVIDIA A100

The Ampere-generation specs (table E.4) anchor most training examples in the book.

Table E.4: NVIDIA A100 (Ampere): Peak specs from NVIDIA’s architecture whitepaper, 80 GB datasheet, and architecture article (NVIDIA Corporation 2020a, 2021a; Choquette et al. 2021). The A100 is the most commonly cited accelerator in training examples; 80 GB HBM2e and TF32 Tensor Cores set memory-capacity and compute-intensity anchors.

Assumption	Value	Unit
Peak FP16 tensor throughput (A100)	312	TFLOP/s
Peak FP32 throughput (A100)	19.5	TFLOP/s
Peak INT8 throughput (A100)	624	TOPS
Peak TF32 throughput (A100)	156	TFLOP/s
HBM bandwidth (A100)	2039	GB/s
HBM capacity (A100)	80	GB
TDP (A100)	400	W

NVIDIA H100

For Hopper-era estimates, the reference values include FP8 Tensor Cores and the Transformer Engine (table E.5).

Table E.5: NVIDIA H100 (Hopper): Peak specs from the Hopper architecture overview (Choquette 2023). FP8 Tensor Cores and the Transformer Engine drive Hopper-generation training and serving estimates.

Assumption	Value	Unit
Peak FP16 tensor throughput (H100)	989	TFLOP/s
Peak FP32 CUDA throughput (H100)	67	TFLOP/s
Peak FP8 tensor throughput (H100)	1979	TFLOP/s
Peak INT8 throughput (H100)	1979	TOPS
Peak TF32 throughput (H100)	494	TFLOP/s
HBM bandwidth (H100)	3.35	TB/s
HBM capacity (H100)	80	GB
TDP (H100)	700	W

NVIDIA B200

Forward-looking capacity-planning examples rely on Blackwell-generation specs (table E.6).

Table E.6: NVIDIA B200 (Blackwell): Peak specs from NVIDIA Blackwell product documentation (NVIDIA Corporation 2024). Forward-looking capacity-planning examples use its HBM3e bandwidth and FP8 throughput.

Assumption	Value	Unit
Peak FP16 tensor throughput (B200)	2250	TFLOP/s
Peak FP8 tensor throughput (B200)	4500	TFLOP/s
Peak FP4 throughput (B200)	9000	TFLOP/s
HBM bandwidth (B200)	8	TB/s
HBM capacity (B200)	180	GB
TDP (B200)	1000	W

AMD Instinct MI300X

To compare with H100-class hardware used elsewhere, MI300X capacity, bandwidth, throughput, and power specifications provide a cross-vendor baseline (table E.7).

Table E.7: AMD Instinct MI300X: Peak specs from AMD Instinct MI300X product documentation (AMD 2023). High HBM capacity (192 GB) and bandwidth make it a common cross-vendor baseline for memory-bound workloads.

Assumption	Value	Unit
Peak FP16 tensor throughput (MI300X)	1307	TFLOP/s
HBM bandwidth (MI300X)	5.3	TB/s
HBM capacity (MI300X)	192	GB
TDP (MI300X)	750	W

Google TPU v4 and v6e

When comparing training economics across accelerator families, TPU figures offer an application-specific integrated circuit (ASIC)-based alternative (table E.8).

Table E.8: Google TPU v4 and v6e (Trillium): TPU v4 values come from Google TPU publications (Jouppi et al. 2023), and TPU v6e values come from Google Cloud's current Trillium specification page (Google Cloud 2026d).

Assumption	Value	Unit
Peak BF16 throughput (TPU v4)	275	TFLOP/s
Memory bandwidth (TPU v4)	1200	GB/s
Peak BF16 throughput (TPU v6e)	918	TFLOP/s
Memory bandwidth (TPU v6e)	1638	GB/s

CPU and mobile/edge processors

Edge and mobile ML examples rely on a different performance baseline (table E.9), contrasting starkly with data center throughput to illustrate why deployment target shapes every design decision.

Table E.9: CPU, Mobile NPU, and Edge Device Specs: Illustrative reference points for edge and mobile examples (not vendor peaks for a single SKU). The contrast with data-center accelerators shows why deployment target shapes every design decision.

Assumption	Value	Unit
Peak FP32 throughput (reference CPU)	1	TFLOP/s
DRAM bandwidth (reference server)	50	GB/s

Assumption	Value	Unit
Reference NPU throughput (iPhone 15 Pro)	35	TOPS
Memory bandwidth (iPhone 15 Pro)	51.2	GB/s
TDP (mobile device, reference)	5	W
Power (edge object detector, reference)	2	W
Battery capacity (phone, reference)	15	Wh

Model Specifications

These are the **reference model assumptions** behind training-cost, memory-footprint, and inference-workload examples: parameter counts, per-inference FLOP budgets, and published training-scale anchors. Published model sizes follow primary papers, model reports, and official model documentation (Devlin et al. 2019; Radford et al. 2019; Brown et al. 2020; Dubey et al. 2024; He et al. 2016a; Sandler et al. 2018; Ultralytics 2023). The inference FLOP budgets are book-derived values: they use sequence length 128 for BERT-Base, 224×224 images for ResNet-50 and MobileNetV2, and 640×640 images for YOLOv8-Nano; the vision models count two FLOPs per multiply-add. GPT-3 training FLOPs follow (Brown et al. 2020); the 25-day, 1,024-A100 pairing is editorial. The GPT-4 parameter count and training GPU-days are public third-party mixture of experts (MoE) estimates (Patel and Wong 2023) because the GPT-4 technical report does not disclose architecture size. When a chapter estimates GPT-3-scale training time, it uses table E.10.

Table E.10: Reference Model Specifications: Parameter counts and FLOP budgets for worked examples. BERT, GPT-2, GPT-3, Llama, ResNet-50, MobileNetV2, and YOLOv8 rows use primary papers, model reports, or official model documentation (Devlin et al. 2019; Radford et al. 2019; Brown et al. 2020; Dubey et al. 2024; He et al. 2016a; Sandler et al. 2018; Ultralytics 2023). GPT-4 rows cite (OpenAI et al. 2023) for the official model family and (Patel and Wong 2023) for the public MoE parameter and GPU-day estimates used in this edition.

Assumption	Value	Unit
Inference FLOPs (BERT-Base)	2.2e+10	flop
Parameters (BERT-Base)	1.1e+08	param
Parameters (Llama 3 8B)	8.03e+09	param
Hidden dimension (GPT-2)	1600	-
Layers (GPT-2)	48	-
Parameters (GPT-2)	1.5e+09	param
Parameters (GPT-3)	1.75e+11	param
Scenario duration (GPT-3 scale)	25	d
Reference training FLOPs (GPT-3)	3.14e+23	flop
Parameters (GPT-4, public MoE estimate)	1.76e+12	param
Reference training GPU-days (GPT-4)	2.5e+06	GPU-days
Inference FLOPs (ResNet-50)	8.2e+09	flop
Parameters (ResNet-50)	2.56e+07	param
Inference FLOPs (MobileNetV2)	6e+08	flop
Parameters (MobileNetV2)	3.50487e+06	param
Inference FLOPs (YOLOv8-Nano)	8.7e+09	flop

Training Memory Conventions

Training-memory napkin math assumes the mixed-precision Adam storage model used in the napkin-math callout and several training chapters: BF16 weights, BF16 gradients, and FP32 master weights plus Adam first- and second-moment buffers (16 bytes per parameter in total). This is a book convention, not a measured hardware constant. The BF16 data path follows established BFLOAT16 mixed-precision practice (Kalamkar et al. 2019); the FP32-master pattern follows earlier FP16 work (Micikevicius et al. 2017; NVIDIA 2017), and Adam’s first- and second-moment buffers

follow (Kingma and Ba 2015). Table E.11 lists per-component byte widths; multiplying bytes per parameter (mixed-precision Adam) by the parameter count gives the nonactivation training-state footprint.

Table E.11: Training Memory Conventions: Per-parameter storage for mixed-precision Adam ($2 + 2 + 12 = 16$ bytes before activations). Book convention for napkin math; see (NVIDIA 2017) for mixed-precision training context.

Assumption	Value	Unit
Weight/gradient width (BF16)	2	bytes
Master weight width (FP32)	4	bytes
Master + Adam states per parameter (FP32)	12	bytes
Bytes per parameter (mixed-precision Adam)	16	bytes

Hardware and model assumptions fix *what* runs where; **energy assumptions** fix whether the design is thermally and economically viable at the operation and memory-access level.

Energy Constants

Horowitz’s 45 nm arithmetic and 32-bit DRAM estimates (Horowitz 2014), together with illustrative book anchors for register and SRAM access, MobileNetV2 inference, and 5G transfer, underpin the book’s efficiency comparisons but are not process-independent. Table E.12 lists the access-energy hierarchy from registers through DRAM. That hierarchy quantifies why data reuse dominates kernel design.

Table E.12: Energy per Operation and Access: Arithmetic and 32-bit DRAM estimates use Horowitz’s 45 nm model (Horowitz 2014), and the per-byte DRAM row is derived from that estimate. Register, SRAM, MobileNetV2, and 5G rows are illustrative book anchors. The roughly $6,400\times$ gap between a 0.1 pJ register access and a 640 pJ DRAM access explains why data reuse dominates ML kernel optimization.

Assumption	Value	Unit
Register access energy	0.1	pJ
L1 SRAM access energy	0.5	pJ
L2 SRAM access energy	2	pJ
DRAM access energy (32-bit)	640	pJ
DRAM access energy (per byte)	160	pJ/byte
FP16 multiply energy (45 nm)	1.1	pJ/multiply
FP32 multiply energy (45 nm)	3.7	pJ/multiply
INT8 multiply energy (45 nm)	0.2	pJ/multiply
MobileNetV2 inference energy (reference)	0.1	mJ
5G transfer energy per MB	100	mJ/MB

Precision & arithmetic energy hierarchy

The arithmetic energy hierarchy in table E.13 compares the energy required per operation (E_{op} in pJ/op) for additions and multiplications across floating-point, tensor, and integer precisions on modern 7nm, 5nm, and 3nm silicon process nodes.

Table E.13: Precision & Arithmetic Energy Hierarchy: Energy per operation (E_{op} in pJ/op) for Addition (Add) and Multiplication (Mul) across data formats (FP64, FP32, TF32, FP16, BF16, FP8 E4M3/E5M2, INT32, INT8, INT4) and semiconductor process nodes (7nm, 5nm, 3nm). Energy scaling illustrates the quadratic $O(b^2)$ multiplier vs linear $O(b)$ adder bitwidth dependence (Horowitz 2014; Dally et al. 2021).

Precision/Data Format	7nm Node (E_{op} pJ)	5nm Node (E_{op} pJ)	3nm Node (E_{op} pJ)	Energy Ratio vs FP32 (3nm)
FP64 (64-bit Float)	Add: 0.90/Mul: 2.50	Add: 0.58/Mul: 1.62	Add: 0.38/Mul: 1.05	$2.23\times$

Precision/Data Format	7nm Node (E_{op} pJ)	5nm Node (E_{op} pJ)	3nm Node (E_{op} pJ)	Energy Ratio vs FP32 (3nm)
FP32 (32-bit Float)	Add: 0.40/Mul: 1.10	Add: 0.26/Mul: 0.72	Add: 0.17/Mul: 0.47	1.00× (Ref)
TF32 (19-bit TensorFloat)	Add: 0.28/Mul: 0.65	Add: 0.18/Mul: 0.42	Add: 0.12/Mul: 0.27	0.57×
FP16 (16-bit IEEE Float)	Add: 0.20/Mul: 0.45	Add: 0.13/Mul: 0.29	Add: 0.08/Mul: 0.19	0.40×
BF16 (16-bit Bfloat16)	Add: 0.18/Mul: 0.40	Add: 0.12/Mul: 0.26	Add: 0.08/Mul: 0.17	0.36×
FP8 (E4M3/E5M2)	Add: 0.09/Mul: 0.20	Add: 0.06/Mul: 0.13	Add: 0.04/Mul: 0.08	0.17×
INT32 (32-bit Integer)	Add: 0.10/Mul: 0.80	Add: 0.06/Mul: 0.52	Add: 0.04/Mul: 0.34	0.72×
INT8 (8-bit Integer)	Add: 0.03/Mul: 0.20	Add: 0.02/Mul: 0.13	Add: 0.013/Mul: 0.085	0.18×
INT4 (4-bit Integer)	Add: 0.010/Mul: 0.060	Add: 0.006/Mul: 0.039	Add: 0.004/Mul: 0.025	0.053×

Energy costs operate at the chip level, but real ML systems also move data across interconnects—between accelerators, across racks, and over wide-area networks. The next section lists the bandwidth assumptions used when chapters estimate communication overhead.

Interconnect and Network Bandwidth

These **bandwidth assumptions** apply when chapters reason about gradient synchronization, pipeline bubbles, checkpoint I/O, or cross-data center latency. NVLink and PCIe rates follow accelerator product documentation (NVIDIA Corporation 2017, 2020a; Choquette 2023); the InfiniBand architecture specification anchors the protocol family (InfiniBand Trade Association 2000), while current high-speed product families and Ethernet roadmaps anchor modern link-rate examples (NVIDIA 2026f; Ethernet Alliance 2025); the speed-of-light-in-fiber floor is a physics identity. Table E.14 lists NVLink, InfiniBand, PCIe, Non-Volatile Memory Express (NVMe), and Ethernet rates together with the propagation speed of light in fiber.

Table E.14: Interconnect and Network Bandwidth: Link-family bandwidths use vendor specs, the InfiniBand architecture specification, and current product/roadmap documentation (InfiniBand Trade Association 2000; NVIDIA 2026f; Ethernet Alliance 2025). NVLink rates are aggregate bidirectional; other link rates are nominal per direction. Speed of light in fiber sets the cross-data-center latency floor.

Assumption	Value	Unit
NVLink bandwidth (V100)	300	GB/s
NVLink bandwidth (A100)	600	GB/s
NVLink bandwidth (H100)	900	GB/s
InfiniBand HDR link rate	200	Gb/s
InfiniBand NDR link rate	400	Gb/s
InfiniBand XDR link rate	800	Gb/s
PCIe Gen4 ×16 rate	32	GB/s
PCIe Gen5 ×16 rate	64	GB/s
NVMe sequential read bandwidth	7	GB/s
10 GbE link rate	10	Gb/s
100 GbE link rate	100	Gb/s
Speed of light in fiber	200000	km/s

Interconnect & network fabric physical spec sheet

The physical specification sheet in table E.15 details bandwidth, latency, framing, direct memory access (DMA) support, and energy per bit across host buses, accelerator interconnects, network fabrics, and cache-coherent interfaces.

Table E.15: Interconnect & Network Fabric Physical Spec Sheet: Comprehensive physical spec sheet for host buses, GPU interconnects, switches, network fabrics, and coherent interfaces (PCIe Gen3–Gen7, NVLink 1–5, NVSwitch 1–4, InfiniBand EDR–XDR, RoCE v2, CXL 1.1–3.1). Columns enumerate directional and aggregate bandwidth, PHY/link latency, line encoding/framing, DMA and GPUDirect capability, and physical layer energy cost (E_{link} in pJ/bit).

Bus/Interconnect Standard	Directional & Aggregate Bandwidth	Latency	Encoding/Framing	DMA/GPUDirect Support	Energy Cost (E_{link})
PCIe Gen3 (x16)	15.75 GB/s dir/31.5 GB/s agg	400–500 ns	128b/130b NRZ	PCIe P2P, GPUDirect RDMA	14.0 pJ/bit
PCIe Gen4 (x16)	31.50 GB/s dir/63.0 GB/s agg	250–400 ns	128b/130b NRZ	PCIe P2P, GPUDirect RDMA	10.0 pJ/bit
PCIe Gen5 (x16)	63.00 GB/s dir/126.0 GB/s agg	150–250 ns	128b/130b NRZ	PCIe P2P, GPUDirect RDMA, CXL 1.1/2.0	8.0 pJ/bit
PCIe Gen6 (x16)	126.00 GB/s dir/252.0 GB/s agg	100–150 ns	PAM4 (256B Flit)	PCIe P2P, GPUDirect RDMA, CXL 3.0/3.1	6.0 pJ/bit
PCIe Gen7 (x16)	252.00 GB/s dir/504.0 GB/s agg	<100 ns	PAM4 (512B Flit)	PCIe P2P, GPUDirect RDMA, CXL 3.1+	4.5 pJ/bit
NVLink 1 (P100)	20 GB/s dir/40 GB/s agg (link)	200–300 ns	NRZ (20 Gb/s)	GPUDirect P2P/RDMA	10.0 pJ/bit
NVLink 2 (V100)	25 GB/s dir/50 GB/s agg (link)	150–200 ns	NRZ (25.78 Gb/s)	GPUDirect P2P/RDMA	8.0 pJ/bit
NVLink 3 (A100)	25 GB/s dir/50 GB/s agg (link)	100–150 ns	NRZ (50 Gb/s)	GPUDirect P2P/RDMA	6.5 pJ/bit
NVLink 4 (H100)	25 GB/s dir/50 GB/s agg (link)	80–100 ns	PAM4 (100 Gb/s)	GPUDirect P2P, SHARP Aggregation	4.5 pJ/bit
NVLink 5 (B200)	50 GB/s dir/100 GB/s agg (link)	60–80 ns	PAM4 (200 Gb/s)	GPUDirect P2P, NVLink Network Offload	3.5 pJ/bit
NVSwitch 1 (Volta)	900 GB/s aggregate per chip	~100 ns	NRZ	Hardware P2P Crossbar Routing	8.5 pJ/bit
NVSwitch 2 (Ampere)	2.4 TB/s aggregate per chip	~90 ns	NRZ	Hardware P2P Crossbar, SHARP v2	6.5 pJ/bit
NVSwitch 3 (Hopper)	3.2 TB/s aggregate per chip	~70 ns	PAM4	SHARP v3 In-Network Compute	4.5 pJ/bit
NVSwitch 4 (Blackwell)	14.4 TB/s aggregate per chip	~50 ns	PAM4	SHARP v4, FP8 Reduction Engine	3.2 pJ/bit
InfiniBand EDR	100 Gbps (12.5 GB/s dir)	~0.50 μ s (500 ns)	64b/66b NRZ	GPUDirect RDMA (Verbs)	15.0 pJ/bit
InfiniBand HDR	200 Gbps (25.0 GB/s dir)	~0.60 μ s (600 ns)	64b/66b PAM4/NRZ	GPUDirect RDMA, SHARP v2	10.0 pJ/bit
InfiniBand NDR	400 Gbps (50.0 GB/s dir)	~0.50 μ s (500 ns)	256b/257b PAM4	GPUDirect RDMA, SHARP v3	7.0 pJ/bit
InfiniBand XDR	800 Gbps (100.0 GB/s dir)	~0.40 μ s (400 ns)	PAM4 (200G/lane Flit)	GPUDirect RDMA, SHARP v4	5.0 pJ/bit
RoCE v2 Ethernet	100–800 Gbps (12.5–100 GB/s dir)	1.0–2.5 μ s	NRZ/PAM4 (802.3)	GPUDirect RDMA (RoCEv2 UDP/IP)	8.0–12.0 pJ/bit
CXL 1.1	63 GB/s dir/126 GB/s agg (x16)	~200 ns	128b/130b NRZ	cxl.io, cxl.cache, cxl.mem	8.0 pJ/bit
CXL 2.0	63 GB/s dir/126 GB/s agg (x16)	~180 ns	128b/130b NRZ	SLD/MLD Memory Pooling	7.5 pJ/bit
CXL 3.0	126 GB/s dir/252 GB/s agg (x16)	~120 ns	256B Flit PAM4	P2P Memory Pooling & Fabric Switch	5.5 pJ/bit
CXL 3.1	126 GB/s dir/252 GB/s agg (x16)	~110 ns	256B Flit PAM4	Global Integrated Memory, TEE	5.0 pJ/bit

Economic Constants

These **pricing assumptions** underpin total cost of ownership (TCO) and energy-cost napkin math in table E.16. They are illustrative hyperscaler-order rates for ratio analysis (similar in spirit to carbon

accounting examples in (Patterson et al. 2021)), not quotes for a specific region or contract. Local values should replace these anchors when absolute price dominates.

Table E.16: Economic Assumptions: Illustrative cloud electricity and egress rates for TCO napkin math (2024–2025 order of magnitude). On-premise cost structures differ; relative magnitudes guide design trade-offs.

Assumption	Value	Unit
Cloud electricity price	0.12	dollar/kWh
Cloud egress price per GB	0.09	dollar/GB

Economic constants set the price per unit of compute and data transfer, but they mean little without a sense of the volumes involved. Production ML systems handle millions to billions of requests per day—numbers large enough to be difficult to internalize without concrete reference points.

Scale References

These **scale assumptions** in table E.17 anchor “how big is big?” for capacity-planning examples, providing illustrative order-of-magnitude baselines for email and search volume, autonomous-vehicle sensor streams, and standard 1080p and 4K video formats. They are magnitude anchors, not audited statistics for a specific year.

Table E.17: Production Scale and Data Rate References: Illustrative workload and sensor-rate anchors, plus standard 1080p and UHD 4K video parameters.

Assumption	Value	Unit
Gmail emails per day	1.21e+11	-
Google searches per day	8.5e+09	-
Waymo sensor data rate (low)	1	TB/h
Waymo sensor data rate (high)	19	TB/h
1080p frame width	1920	-
1080p frame height	1080	-
4K frame width	3840	-
4K frame height	2160	-
Bytes per RGB pixel	3	bytes
Video frame rate (standard)	30	Hz

The assumptions above use shared unit conventions: decimal data prefixes, distinct FLOPs and FLOP/s quantities, and the aliases defined in the next section.

Unit Conventions

Table E.18 fixes the unit conventions used in every quantitative example in this book. Each row gives the multiplier k in $1 \text{ alias} = k \text{ base}$. Data prefixes use decimal SI ($\text{KB} = 10^3$ bytes, not 1024). Binary IEC storage prefixes appear in a few storage-specific discussions but are omitted here because most fleet-scale estimates in the book use decimal KB/GB/TB. Throughput quantities (FLOP/s, GB/s) combine these aliases with time; adding incompatible dimensions (bytes to FLOP/s) is a category error in napkin math, not a unit conversion.

Hardware capacity is the one deliberate exception. Vendors label memory with decimal symbols but ship binary quantities. An accelerator sold as an 80 GB device provides 80×2^{30} bytes, and a microcontroller sold with 512 KB of SRAM provides 512×2^{10} bytes. This book keeps both the vendor’s number and the vendor’s label, so a capacity printed as 80 GB is the nameplate figure rather than a decimal conversion of it. Every other data quantity, including model footprints, activation sizes, and transfer volumes, uses the decimal prefixes above. Where a worked example subtracts a

computed footprint from a device capacity, both terms are carried in decimal so that the arithmetic printed on the page is the arithmetic a reader can reproduce.

Table E.18: Unit Conventions: Decimal SI aliases and scale factors (book convention for napkin math). Throughput forms such as FLOP/s and GB/s divide work or data by time using the same conventions.

Alias	Multiplier	Base unit
byte	1	byte
KB	1000	byte
MB	1e+06	byte
GB	1e+09	byte
TB	1e+12	byte
PB	1e+15	byte
flop	1	flop
GFLOPs	1e+09	flop
TFLOPs	1e+12	flop
ZFLOPs	1e+21	flop
param	1	param
Mparam	1e+06	param
Gbps	1e+09	bit/s
NS	10 ^{−9}	second
US	10 ^{−6}	second
MS	10 ^{−3}	second
second	1	second
hour	3600	second
day	86400	second
joule	1	joule
watt	1	watt
meter	1	meter

Each assumption also has a provenance record that lets readers trace the number back to its primary source.

E.1 Assumption Provenance

Table E.19 groups the appendix’s quantitative inputs by evidence class and identifies the primary references behind each group. Vendor specifications and standards define published limits, peer-reviewed measurements provide empirical anchors, and explicitly marked editorial scenarios supply illustrative values where no durable universal constant exists. This distinction lets readers judge the evidentiary weight of an input, update time-sensitive values, and reproduce an estimate with an appropriate local substitute.

Table E.19: Assumption Provenance Catalog: Quick map from appendix section to source class and bibliography.

Appendix section	Source type	Primary references
Accelerator specifications	Vendor datasheet peaks; editorial CPU/mobile anchors; Master Spec Matrix (table E.1)	(NVIDIA Corporation 2017, 2018, 2020a, 2024; Choquette et al. 2021; Choquette 2023; AMD 2023; Jouppi et al. 2023; Google Cloud 2026d); editorial for CPU/mobile rows
Model specifications	Published papers, model reports, and official docs; editorial GPT-3 scenario; GPT-4 size from public analysis	(Devlin et al. 2019; Radford et al. 2019; Brown et al. 2020; Dubey et al. 2024; He et al. 2016a; Sandler et al. 2018; Ultralytics 2023; OpenAI et al. 2023; Patel and Wong 2023)

Appendix section	Source type	Primary references
Training memory conventions	Book convention (mixed-precision Adam layout)	(Kingma and Ba 2015; Kalamkar et al. 2019; Micikevicius et al. 2017; NVIDIA 2017)
Energy constants	Published 45 nm arithmetic/DRAM estimates; VLSI synthesis node scaling; Precision Energy Hierarchy (table E.13)	(Horowitz 2014; Dally et al. 2021) for arithmetic/DRAM; editorial for register/SRAM, MobileNetV2, and 5G rows
Interconnect bandwidth	Vendor specs; InfiniBand standard; PCIe/NVLink/CXL standards; Physical Spec Sheet (table E.15)	(NVIDIA Corporation 2017, 2020a; Choquette 2023; InfiniBand Trade Association 2000; NVIDIA 2026f; Ethernet Alliance 2025)
Economic assumptions	Illustrative cloud/utility rates	(Patterson et al. 2021) (methodology context)
Scale references	Illustrative workload rates; standard video conventions	Editorial for workload rates; standard format definitions for video
Unit conventions	Decimal SI; book notation	Editorial

Summary

This appendix is the reference catalog for the quantitative inputs used throughout the book. It gathers accelerator and model specifications, training-memory conventions, energy constants, interconnect bandwidths, economic and scale references, unit conventions, and provenance. The values come from vendor peaks, published measurements, standards, and editorial scenarios, so they carry different evidentiary weight and age at different rates.

Use these tables as starting inputs, not universal constants. For each estimate, record the table row, unit convention, and provenance; replace time-sensitive peaks, prices, and rates with local values; distinguish a physical ceiling from a measured operating point and a scenario assumption; and report which assumptions dominate the result. This keeps napkin math reproducible and makes clear which input controls the conclusion.

Glossary

This glossary defines key terms used throughout this book and gathers them as a compact reference. Terms are organized alphabetically so they can be consulted alongside the chapters.

3

3DMark Graphics performance benchmark suite that evaluates real-time 3D rendering capabilities, measuring triangle throughput, texture fill rates, and modern features like ray tracing and DLSS performance.

A

A/B testing A controlled experimental method for comparing two versions of a system or model by randomly dividing users into groups and measuring performance differences between the variants.

ablation study An experiment that removes, disables, or isolates one component at a time to measure its contribution to model or system behavior.

activation-based pruning A pruning method that evaluates the average activation values of neurons or filters over a dataset to identify and remove neurons that consistently produce low activations and contribute little information to the network's decision process.

activation checkpointing A memory optimization technique that reduces memory usage during backpropagation by selectively discarding and recomputing activations instead of storing all intermediate results.

activation function A mathematical function applied to the weighted sum of inputs in a neural network neuron to introduce non-linearity, enabling the network to learn complex patterns beyond simple linear combinations.

activation memory Memory used to store intermediate layer outputs needed for gradient computation during training, scaling with batch size, sequence length, and model depth.

active learning An approach that intelligently selects the most informative examples for human annotation based on model uncertainty, reducing the amount of labeled data needed for effective training.

Adam optimization An adaptive learning rate optimization algorithm that combines momentum and RMSprop by maintaining exponentially decaying averages of both gradients and squared gradients for each parameter.

AdamW A variant of Adam that decouples weight decay from the adaptive gradient update, improving regularization while keeping the optimizer's adaptive moments.

AI Triad A framework modeling ML systems as three interdependent components: data that guides behavior, algorithms that learn pat-

terns, and computational infrastructure that enables training and inference. Limitations in any component constrain the capabilities of the others.

alerting Automated notification systems that inform teams when metrics exceed predefined thresholds or anomalies are detected in production ML systems.

AlexNet A landmark convolutional neural network architecture that won the 2012 ImageNet challenge, reducing error rates from 26 percent to 16 percent and sparking the deep learning renaissance.

AllReduce A collective communication operation in distributed computing where each process contributes data and all processes receive the combined result, commonly used for gradient aggregation in distributed training.

all-to-all communication A communication pattern where every device or process exchanges data with every other device, common in sharded embedding and expert-parallel workloads.

AlphaFold A landmark AI system developed by DeepMind that predicts the three-dimensional structure of many proteins from their amino acid sequences, achieving breakthrough accuracy without solving every folding problem and demonstrating how large-scale ML systems can accelerate scientific discovery. Its predictions still require confidence interpretation and, for many uses, experimental validation.

Amdahl's Law A speedup bound showing that overall improvement is limited by the portion of a workload or pipeline that remains serial or unoptimized.

AOT compilation Ahead-of-time compilation that converts a model, graph, or program into optimized executable code before deployment or execution, trading flexibility for lower runtime overhead.

Apache Kafka A distributed streaming platform that handles real-time data feeds using a publish-subscribe messaging system, commonly used for building ML data pipelines with high throughput and fault tolerance.

Apache Spark An open-source distributed computing framework that enables large-scale data processing across clusters of computers, revolutionizing ETL operations with in-memory computing capabilities.

application-specific integrated circuit (ASIC) Custom chips designed for specific computational tasks that offer superior performance and energy efficiency compared to general-purpose processors by abandoning general-purpose flexibility. Examples include Google's TPUs, Cerebras Wafer-Scale Engine, and Bitcoin mining ASICs.

architectural efficiency The dimension of model optimization that focuses on how computations are performed efficiently during training and inference by exploiting sparsity, factorizing large components, and dynamically adjusting computation based on input complexity.

arithmetic intensity The ratio of floating-point operations to bytes of memory accessed (FLOP/byte), used in roofline analysis to determine whether workloads are memory bound or compute bound and to guide optimization priorities.

artificial general intelligence AI systems intended to match human-level performance across a broad range of cognitive tasks; the term has no universally accepted capability threshold and implies no particular architecture.

artificial intelligence The field of computer science focused on creating systems that can perform tasks typically requiring human intelligence, such as perception, reasoning, learning, and decision-making.

artificial neural network A computational model inspired by biological neural networks, consisting of interconnected nodes (neurons) organized in layers that can learn patterns from data through adjustable weights and biases.

artificial neurons Basic computational units in neural networks that mimic biological neurons, taking multiple inputs, applying weights and biases, and producing an output signal through an activation function.

attention mechanism A neural network component that computes weighted connections between elements based on their content, allowing dynamic focus on relevant parts of the input rather than fixed architectural connections.

AUC Area under the ROC curve, a threshold-independent classification metric that compares true positive and false positive rates across decision thresholds.

audit trail An append-only record of who accessed or changed data, models, or system artifacts and when, supporting accountability, debugging, and compliance.

autograd tape A transient record of operations and saved values built during forward execution so reverse-mode automatic differentiation can compute gradients during the backward pass.

automatic differentiation A computational technique that automatically calculates exact derivatives of functions implemented as computer programs by systematically applying the chain rule at the elementary operation level, essential for training

neural networks through gradient-based optimization.

automatic mixed precision A training technique that automatically manages the use of different numerical precisions (FP16, FP32) to optimize memory usage and computational speed while maintaining model accuracy.

AutoML Automated machine learning that automates model design decisions, including architecture search, hyperparameter optimization, and feature selection, to reduce manual effort.

autoscaling Dynamic adjustment of compute resources based on workload demand, automatically scaling up during peak usage and scaling down during low usage to optimize costs and performance.

B

backpropagation An algorithm that computes gradients of the loss function with respect to network weights by propagating error signals backward through the network layers, enabling systematic weight updates during training.

backward pass The training phase that propagates gradients from the loss back through the network to compute parameter gradients, pairing with the forward pass.

bandwidth The maximum rate of data transfer across a communication channel or memory interface, typically measured in bytes per second and critical for optimizing data movement in AI accelerators.

batch inference The process of running inference on a large dataset in bulk, typically as a scheduled job, as opposed to real-time inference on individual requests. Batch inference enables higher throughput by amortizing overhead across many inputs.

batch ingestion A data processing pattern that collects and processes data in groups or batches at scheduled intervals, suitable for scenarios where real-time processing is not critical.

batch normalization A technique that normalizes inputs to each layer to have zero mean and unit variance, which stabilizes training and often allows for higher learning rates and faster convergence, and subsequently applies a learnable scale and shift to preserve representational power.

batch processing The technique of processing multiple data samples simultaneously to amortize computation and memory access costs, improving overall throughput in neural network training and inference.

batch size The number of training examples processed simultaneously during one iteration of neural network training, affecting both computational efficiency and gradient estimation quality.

batch throughput optimization Techniques for maximizing the number of samples processed per unit time when handling multiple inputs simultaneously, using parallelism and batching efficiencies.

batched operations Matrix computations that process multiple inputs simultaneously, converting matrix-vector operations into more efficient matrix-matrix operations to improve hardware utilization.

benchmark engineering The systematic design and development of performance evaluation frameworks, involving test harness creation, metric selection, and result interpretation methodologies.

benchmark harness Systematic infrastructure component that controls test execution, manages input delivery, and collects perfor-

mance measurements under controlled conditions to ensure reproducible evaluations.

benchmarking Systematic evaluation of compute performance, algorithmic effectiveness, and data quality in machine learning systems to optimize performance across diverse workloads and ensure reproducibility.

BERT Bidirectional Encoder Representations from Transformers, a transformer-based language model introduced by Google in 2018 that revolutionized natural language processing through masked language modeling pretraining.

BF16 A 16-bit floating-point format developed by Google Brain that maintains the same dynamic range as FP32 but with reduced precision, making it particularly suitable for deep learning training.

bias A learnable parameter added to the weighted sum in each neuron that shifts the activation function, allowing neurons to activate even when all inputs are zero and providing additional flexibility for the network to fit complex patterns.

binarization An extreme quantization technique that reduces neural network weights, activations, or both to binary values (typically -1 and +1), achieving maximum compression but often requiring specialized training procedures and hardware support.

biological neuron A cell in the nervous system that receives, processes, and transmits information through electrical and chemical signals, serving as inspiration for artificial neural networks.

bitter lesson Richard Sutton's 2019 observation that general methods using computation consistently outperform approaches encoding human expertise, suggesting that systems engineering enabling computational scale is central to AI advancement.

black box A system whose inputs and outputs are observable but whose internal workings are inaccessible or opaque. This opacity is particularly problematic when the system makes consequential decisions without explaining its reasoning.

BLAS Basic Linear Algebra Subprograms, a specification for low-level routines that perform common linear algebra operations such as vector addition, scalar multiplication, dot products, and matrix operations, forming the computational foundation of modern ML frameworks.

block sparse matrix A sparse matrix whose nonzero values appear in dense blocks rather than isolated elements, preserving structure that hardware kernels can exploit efficiently.

blue-green deployment A deployment strategy using two production environments so traffic can switch to a validated new version while the old version remains available for rollback.

bounding box A rectangular annotation that identifies object locations in images by drawing a box around each object of interest, commonly used in computer vision training datasets.

brittleness The tendency of rule-based AI systems to fail completely when encountering inputs that fall outside their programmed scenarios, no matter how similar those inputs might be to what they were designed to handle.

C

caching A technique for storing frequently accessed data in high-speed storage systems to

reduce retrieval latency and improve system performance in ML pipelines.

caching allocator A framework memory manager that reuses device-memory blocks instead of calling the hardware allocator for every tensor, reducing allocation latency and fragmentation.

calibration The process in post-training quantization of analyzing a representative dataset to determine optimal quantization parameters, including scale factors and zero points, that minimize accuracy loss when converting from high to low precision.

canary deployment A deployment strategy where new model versions receive a small percentage of production traffic to validate behavior before full rollout, enabling early detection of issues with minimal user impact.

CAP theorem A distributed systems principle stating that, during a network partition, a data store cannot guarantee both linearizable consistency and availability for every request.

carbon footprint The total greenhouse gas emissions, typically measured in CO₂ equivalent, produced directly and indirectly by training and operating an ML system.

carbon intensity The emissions produced per unit of energy, usually measured as kg CO₂e per kWh, used to convert ML energy use into carbon impact.

Cerebras Wafer-Scale Engine In the CS-2 generation, a single-wafer processor containing 2.6 trillion transistors and 850,000 cores, designed to eliminate inter-device communication bottlenecks in large-scale machine learning training.

chain rule The calculus rule enabling computation of derivatives for composite functions, fundamental to backpropagation.

channelwise quantization A quantization granularity approach where each channel in a layer uses its own set of quantization parameters, providing more precise representation than layerwise quantization while maintaining hardware efficiency.

CI/CD pipelines Continuous Integration and Continuous Delivery automated workflows that streamline model development by integrating testing, validation, and deployment processes.

circuit breaker pattern A resilience pattern that stops or redirects calls to an unhealthy service after repeated failures to prevent cascading failure and enable fallback behavior.

classification labels Simple categorical annotations that assign specific tags or categories to data examples, representing the most basic form of supervised learning annotation.

CloudSuite Benchmark suite developed at EPFL that addresses modern data center workloads including web search, data analytics, and media streaming, measuring end-to-end performance across network, storage, and compute dimensions.

cold-start performance Time required for a system to transition from idle state to active execution, particularly important in serverless environments where models are loaded on demand.

collaborative filtering A technique used in recommendation systems that predicts user preferences by identifying patterns in interactions from many users, using the collective behavior of the crowd rather than just item properties.

communication tax The performance penalty incurred in distributed systems due to the latency and bandwidth costs of synchronizing state between nodes.

- compound AI systems** AI architectures that combine multiple specialized models and components working together, rather than relying on a single monolithic model, enabling modularity, specialization, and improved interpretability.
- compressed sparse row (CSR)** A memory-efficient storage format for sparse matrices that uses three arrays (values, column indices, row pointers) to store only nonzero elements, reducing memory from $\mathcal{O}(N^2)$ (for an $N \times N$ matrix) to $\mathcal{O}(K + N)$ where K is the number of nonzeros.
- computational graph** A directed graph representing the sequence of operations in neural network computation, enabling automatic differentiation.
- computer engineering** An engineering discipline that emerged in the late 1960s to address the growing complexity of integrating hardware and software systems, combining expertise from electrical engineering and computer science to design and build complex computing systems.
- concept drift** The phenomenon where $P(Y|X)$, the statistical relationship between inputs and targets, changes over time, distinct from data drift in $P(X)$; this can, but need not, degrade a fixed model's performance. Detecting it typically requires outcome labels because input drift alone does not establish a conditional change.
- conditional computation** A dynamic optimization technique where different parts of a neural network are selectively activated based on input characteristics, reducing computational load by skipping unnecessary computations for specific inputs.
- consensus labeling** A quality control approach that collects multiple annotations for the same data point to identify controversial cases and improve label reliability through inter-annotator agreement.
- containerization** Packaging applications and their dependencies into portable, isolated containers using tools like Docker to ensure consistent execution across different environments.
- continuous batching** An LLM serving technique that adds and removes requests from an inference batch between token-generation steps as sequences start and finish.
- continuous integration** A software development practice where code changes are automatically integrated, tested, and validated multiple times per day to detect issues early in the development cycle.
- convolution** A mathematical operation that slides a filter (kernel) across input data to extract features such as edges, textures, or patterns. Convolution is fundamental to convolutional neural networks and particularly effective for processing images and spatial data.
- convolutional neural network** A specialized neural network architecture designed for processing grid-like data such as images, using convolutional layers that apply filters to detect local features.
- coreset** A small subset of a dataset selected to preserve important statistical or training properties of the full dataset while reducing training or evaluation cost.
- correction cascade** An ML technical-debt pattern where fixing one component creates new failures or redesign needs elsewhere because data and model dependencies propagate through the system.
- covariate shift** A distribution shift where input feature distributions change while the relationship between features and labels remains stable.
- CP decomposition** CANDECOMP/PARAFAC decomposition that expresses a tensor as a sum of rank-one components, used to compress neural network layers by reducing the number of parameters while preserving computational functionality.
- credit assignment problem** The challenge of determining which weights in a multi-layer network contributed to prediction errors.
- CRISP-DM** Cross-Industry Standard Process for Data Mining, a structured methodology developed in 1996 that defines six phases for data projects: business understanding, data understanding, data preparation, modeling, evaluation, and deployment.
- critical batch size** The batch size beyond which larger batches provide diminishing optimization benefit and may reduce generalization, limiting simple linear learning-rate scaling.
- cross-entropy loss** A loss function commonly used in classification tasks that measures the difference between predicted probability distributions and true class labels, providing strong gradients for effective learning.
- crowdsourcing** A collaborative data collection approach that uses distributed individuals via the internet to perform annotation tasks, enabling scalable dataset creation through platforms like Amazon Mechanical Turk.
- cuBLAS** NVIDIA's CUDA Basic Linear Algebra Subprograms library that provides GPU-accelerated implementations of standard linear algebra operations, enabling high-performance matrix computations on NVIDIA graphics processing units.
- CUDA (Compute Unified Device Architecture)** NVIDIA's parallel computing platform and programming model that enables general-purpose computing on graphics processing units (GPUs), allowing machine learning frameworks to use massive parallelism for accelerated tensor operations.
- CUDA stream** An ordered queue of GPU work used to schedule kernels, memory transfers, and synchronization events so independent operations can overlap.
- D**
- D-A-M taxonomy** A diagnostic framework that analyzes ML system bottlenecks through three interacting axes: Data (information flow, bounded by bandwidth), Algorithm (mathematical logic, bounded by total operations), and Machine (physical execution, bounded by peak throughput).
- Dartmouth conference** The legendary 8-week workshop at Dartmouth College in 1956 where AI was officially born, organized by John McCarthy, Marvin Minsky, Nathaniel Rochester, and Claude Shannon, following their 1955 proposal, which introduced the term artificial intelligence.
- data as source code** The Software 2.0 principle that training data defines ML system behavior, so dataset changes act like behavior-changing code changes.
- data augmentation** The process of artificially expanding training datasets by creating modified versions of existing data through transformations like rotation, scaling, or noise injection.
- data cascades** Systemic failures where data quality issues compound over time, creating downstream negative consequences such as model failures, costly rebuilding, or project termination.
- data center** A facility that houses computer systems and associated components such as telecommunications and storage systems, typically containing thousands of servers for cloud computing operations.
- data-centric approach** A machine learning paradigm that prioritizes improving data quality, diversity, and curation rather than solely focusing on model architecture improvements to achieve better performance.
- data-centric computing** Systems optimized for the efficient ingestion of data and iterative refinement of model parameters, where the programmer's job is to curate data.
- data contracts** Explicit agreements between data producers and consumers defining schema, quality expectations, and service levels to prevent downstream breakage.
- data curation** The process of selecting, organizing, and maintaining high-quality datasets by removing irrelevant information, correcting errors, and ensuring data meets specific standards for machine learning applications.
- data debt** Accumulated data quality, documentation, schema, lineage, or freshness liabilities that compound over time and degrade model behavior.
- data drift** The phenomenon where the input distribution $P(X)$ changes over time even though code remains unchanged; this may affect model performance but does not guarantee degradation. It can be monitored without labels, although labels are needed to measure its effect on quality.
- data echoing** A training input-pipeline technique that reuses batches, often with fresh augmentation, when data preparation is slower than accelerator consumption.
- data governance** The framework of policies, procedures, and technologies that ensure data security, privacy, compliance, and ethical use throughout the machine learning pipeline.
- data gravity** The architectural pull created by large data volumes, transfer time, bandwidth limits, and egress cost, which tends to draw computation toward the data.
- data ingestion** The process of collecting and importing raw data from various sources into a system where it can be stored, processed, and prepared for machine learning applications.
- data lake** A storage repository that holds structured, semi-structured, and unstructured data in its native format, using schema-on-read approaches for flexible data analysis.
- data lakehouse** A storage architecture that combines data lake flexibility with warehouse-style query semantics, schema management, and transactional guarantees.
- data lineage** The documentation and tracking of data flow through various transformations and processes, providing visibility into data origins and modifications for compliance and debugging.
- data locality** The principle that data and computation should be colocated when transfer latency, bandwidth, or cost dominates remote processing.
- data parallelism** A distributed training strategy that splits the dataset across multiple devices while each device maintains a complete copy of the model, enabling parallel computation of gradients.
- data pipeline** The infrastructure and workflows that automate the movement and transformation of data from sources through processing stages to final storage or consumption.
- data quality** The degree to which data meets requirements for accuracy, completeness, consistency, and timeliness, directly impacting machine learning model performance.

data quality multiplier The idea that improvements in data quality can amplify gains from model and system changes because all downstream stages depend on the data.

data validation The systematic verification that collected data meets quality standards, is properly formatted, and contains accurate information suitable for machine learning model training and evaluation.

data versioning The practice of tracking and managing different versions of datasets over time, similar to code versioning, to ensure reproducibility and enable rollback to previous data states when needed.

data warehouse A centralized repository optimized for analytical queries (OLAP) that stores integrated, structured data from multiple sources in a standardized schema.

dataflow architecture Specialized computing architecture where instruction execution is determined by data availability rather than a program counter, enabling highly parallel processing of neural network operations.

dataflow challenges Technical difficulties in managing data movement and dependencies in hardware accelerators, including memory bandwidth limitations and synchronization requirements.

datasheets for datasets Documentation for training data that captures provenance, collection methodology, demographic composition, and known limitations affecting model behavior.

dead letter queue A separate storage mechanism for data that fails processing, allowing for later analysis and potential reprocessing of problematic data without blocking the main pipeline.

deduplication The removal of exact, near-duplicate, or semantically redundant samples from a dataset to reduce storage, training compute, and leakage risk.

deep learning A subfield of machine learning that uses artificial neural networks with multiple layers to automatically learn hierarchical representations from data without explicit feature engineering.

demographic parity A fairness criterion requiring that the probability of receiving a positive prediction is independent of group membership across protected attributes.

Dennard scaling The observation that as transistors became smaller, their power density remained constant, enabling higher performance at the same power envelope. Its breakdown around 2005 ended the era of free frequency scaling.

dense layer A fully-connected neural network layer where each neuron receives input from all neurons in the previous layer, enabling comprehensive information integration across features.

deployment constraints Operational limitations such as hardware resources, network connectivity, regulatory requirements, and integration requirements that influence how machine learning models are implemented in production environments.

deployment paradigm A distinct approach to hosting and executing ML models characterized by specific resource constraints and operational properties, such as Cloud ML, Edge ML, Mobile ML, or TinyML.

depthwise separable convolution A convolution factorization that applies spatial filtering independently per channel and then mixes channels with a pointwise convolution to reduce computation.

DevOps Software development practice that combines development and operations teams to shorten development cycles and

deliver high-quality software through automation and collaboration.

Dhrystone Integer-based benchmark introduced in 1984 that measures integer and string operations in DMIPS (Dhrystone MIPS), designed to complement floating-point benchmarks with typical programming constructs.

diabetic retinopathy A diabetes complication that damages blood vessels in the retina, serving as a leading cause of preventable blindness and a key application area for medical AI screening systems.

differential privacy A mathematical framework for quantifying and limiting privacy loss when releasing statistical information about datasets by bounding how much the output distribution can change when any one individual's data is added or removed.

direct memory access (DMA) A mechanism that lets a device move data to or from memory without CPU-managed copying for each transfer, enabling overlap and reducing host overhead.

disaggregated evaluation The practice of breaking down model performance metrics by demographic groups or other factors to reveal disparities that are hidden by aggregate measures.

dispatch tax Runtime overhead from launching and orchestrating many small operations through the host framework instead of executing fused or compiled graph regions.

distributed computing An approach that processes data across multiple machines or processors simultaneously, enabling scalable handling of large datasets through frameworks like Apache Spark.

distributed intelligence The placement of computational capabilities across multiple devices and locations rather than relying on a single centralized system, enabling local processing and decision-making.

distributed training A method of training machine learning models across multiple machines or devices to handle larger datasets and models that exceed single-device computational or memory capacity.

distribution shift A change in the statistical properties of data between training and deployment, or over time during deployment. Types include covariate shift (input distribution changes), label shift (output distribution changes), and concept drift (relationship between inputs and outputs changes).

DLRM Deep Learning Recommendation Model, an architecture developed by Meta that combines categorical embeddings with a bottom multi-layer perceptron (MLP) and top MLP to handle the massive scale and sparsity of recommendation system workloads.

domain-specific architecture Hardware designs tailored to optimize specific computational workloads, trading flexibility for improved performance and energy efficiency compared to general-purpose processors.

dropout A regularization technique that randomly sets a fraction of input units to zero during training to prevent overfitting and improve generalization.

dying ReLU problem A failure mode where ReLU neurons become permanently inactive and output zero for all inputs, preventing them from contributing to learning when weighted inputs consistently produce negative preactivations.

dynamic batching A serving technique that waits briefly to group independent requests into a batch, improving throughput while adding controlled batch-formation latency.

dynamic graph A computational graph that is built and modified during program execution, allowing for flexible model architectures and easier debugging but potentially limiting optimization opportunities compared to static graphs.

dynamic pruning A pruning approach that changes which parameters or units are inactive in response to input data or training dynamics, rather than relying only on a fixed pruning decision.

dynamic quantization The process of storing neural-network weights at lower precision and determining activation quantization parameters at runtime, reducing memory use while requiring supported low-precision kernels for speedup.

dynamic random-access memory (DRAM) A type of volatile memory that stores data in capacitors and requires periodic refresh cycles, commonly used as main memory in computer systems.

dynamic voltage and frequency scaling (DVFS) Power management technique that adjusts processor voltage and clock frequency based on workload demands to optimize energy consumption while maintaining performance.

E

eager execution An execution mode where operations are evaluated immediately as they are called in the code, providing intuitive debugging and development experience but potentially sacrificing some optimization opportunities available in graph-based execution.

early exit architectures Neural network designs that include multiple prediction heads at different depths, allowing samples to exit early when confident predictions can be made, reducing average computational cost per inference.

edge computing A distributed computing paradigm that brings computation and data storage closer to the sources of data, reducing latency and bandwidth usage.

edge deployment A deployment strategy where machine learning models run locally on devices at the network edge rather than in centralized cloud servers, reducing latency and enabling operation without constant internet connectivity.

efficiency frontier The optimal trade-off curve between accuracy and computational cost (latency, energy, or memory), where no model exists that is both more accurate and more efficient.

EfficientNet A family of neural network architectures discovered through neural architecture search that achieves better accuracy-efficiency trade-offs by using compound scaling to balance network depth, width, and input resolution.

EL2N (Error L2-Norm) A sample-scoring method that ranks examples by early prediction error, often using a proxy model, to identify informative boundary cases for data selection.

ELIZA One of the first chatbots created by MIT's Joseph Weizenbaum in 1966 that could simulate human conversation through pattern matching and substitution, notable because people began forming emotional attachments to this simple program.

embedded systems Computer systems with dedicated functions within larger mechanical or electrical systems, typically designed for specific tasks with real-time computing constraints.

embedding sharding Splitting large embedding tables across devices or nodes so they fit in aggregate memory, while adding communication to gather embeddings for each batch.

embedding table A lookup table that maps discrete IDs or tokens to dense vectors, often dominating memory in recommendation models and large vocabularies.

emergent behaviors Unexpected system-wide patterns or characteristics that arise from the interaction of individual components, often becoming apparent only when systems operate at scale or in real-world conditions.

encoder-decoder An architectural pattern where an encoder processes input into an intermediate representation and a decoder generates output from this representation, commonly used in sequence-to-sequence tasks.

end-to-end benchmarks Comprehensive evaluation methodology that assesses entire AI system pipelines including data processing, model execution, postprocessing, and infrastructure components.

energy efficiency The measure of computational work performed per unit of energy consumed, typically expressed as operations per joule and crucial for battery-powered and data center deployments.

energy per inference The total energy consumed to produce one model prediction, including model execution and relevant system overhead.

epoch One complete pass through the entire training dataset during neural network training, consisting of multiple batch iterations depending on dataset size and batch size.

equal opportunity A fairness criterion requiring equal true positive rates across different demographic groups.

equalized odds A fairness criterion requiring that both true positive and false positive rates are equal across different demographic groups.

expected calibration error (ECE) A metric measuring the gap between predicted confidence and observed accuracy across confidence bins, used to assess probability calibration.

experiment tracking The systematic recording and management of machine learning experiments, including hyperparameters, model versions, training data, and performance metrics, to enable comparison and reproducibility.

expert systems AI systems first developed in the mid-1960s that captured human expert knowledge in specific domains, exemplified by MYCIN for diagnosing blood infections, representing a shift from general AI to domain-specific applications.

extract, load, transform (ELT) A data processing paradigm that first loads raw data into the target system before applying transformations, providing flexibility for evolving analytical needs.

extract, transform, load (ETL) A traditional data processing paradigm that transforms data before loading it into a data warehouse, resulting in ready-to-query formatted data.

F

fairness laundering A failure mode where removing explicit protected attributes makes a system appear compliant while correlated proxy variables still reproduce discriminatory outcomes.

false positive rate (FPR) The proportion of actual negative cases incorrectly classified as positive, computed as false positives divided by false positives plus true negatives.

FarmBeats A Microsoft Research project that applies machine learning and IoT technologies to agriculture, using edge computing to collect real-time data on soil conditions and crop health while demonstrating distributed AI systems in challenging real-world environments.

feature engineering The process of manually designing and extracting relevant features from raw data to improve machine learning model performance, largely automated in deep learning systems.

feature map The output of a convolutional layer representing the response of learned filters to different spatial locations in the input, capturing detected features at various positions.

feature store A specialized data storage system that provides standardized, reusable features for machine learning, enabling feature sharing across multiple models and teams.

federated learning A machine learning approach that trains algorithms across decentralized edge devices or servers holding local data samples, without exchanging the raw data.

feedback loop A cyclical process where outputs influence future inputs. In ML, this includes both beneficial cycles (where model outputs inform system improvement) and harmful ones (where a model's predictions influence its own future training data, potentially reinforcing and amplifying initial biases over time).

feedforward network A neural network architecture where information flows in one direction from input to output layers without cycles, forming the foundation for many deep learning models.

field-programmable gate array (FPGA) A reconfigurable integrated circuit that can be programmed after manufacturing to implement custom digital circuits and specialized computations, offering flexibility between general-purpose processors and ASICs.

fine-tuning Adapting a pretrained model to a target task or domain by continuing training on task-specific data.

Five-Pillar Framework An organizational structure for ML systems engineering comprising five interconnected disciplines: Data Engineering, Training Systems, Deployment Infrastructure, Operations and Monitoring, and Ethics and Governance.

FlashAttention An IO-aware attention algorithm that tiles computation to avoid materializing the full attention matrix in high-bandwidth memory, reducing memory traffic and enabling longer sequences.

floating-point unit (FPU) A specialized processor component designed to perform arithmetic operations on floating-point numbers with high precision and efficiency.

FLOP/s Floating-Point Operations per Second, a measure of computational throughput that quantifies the number of floating-point operations a system can perform per second.

forgetting events Events where a model changes from classifying an example correctly to misclassifying it later in training, marking difficult or informative samples.

forward pass The process of computing neural network predictions by passing input data through successive layers, applying weights, biases, and activation functions at each stage to produce outputs. Also called *forward propagation*.

foundation model Large-scale machine learning models trained on broad data that can be adapted to a wide range of downstream tasks, serving as a base for specialized applications.

four pillars framework A data engineering framework organizing concerns into Quality, Reliability, Scalability, and Governance to manage the data lifecycle.

FP16 16-bit floating-point numerical representation that reduces memory usage and accelerates computation on modern hardware accelerators while maintaining acceptable precision for many machine learning applications.

FP32 32-bit floating-point numerical representation that provides standard precision for mathematical computations but requires more memory and computational resources than lower-precision formats.

FP32 to INT8 A common quantization transformation that converts 32-bit floating-point weights and activations to 8-bit integers, achieving roughly 4× memory reduction while maintaining acceptable accuracy for many models.

FP8 An 8-bit floating-point format family used on newer accelerators for high-throughput training and inference, requiring careful scaling because of limited precision and range.

framework decomposition The systematic breakdown of neural network frameworks into hardware-mappable components, enabling efficient distribution of operations across processing elements.

G

GEMM General Matrix Multiply operations that follow the pattern $C = \alpha AB + \beta C$, representing the fundamental computational kernel underlying most neural network operations including fully connected layers and convolutional layers.

GEMV General Matrix-Vector multiplication operations that compute the product of a matrix and a vector, commonly used in neural network computations and requiring careful optimization for memory access patterns.

generalization The ability of a machine learning model to perform well on unseen data that differs from the training set, often improved through diverse and high-quality training data.

generalization gap The difference between a model's performance on training data and its performance on held-out data. A large generalization gap can indicate overfitting to the training data.

generative AI A category of artificial intelligence systems capable of creating new content such as text, images, audio, or video based on learned patterns from training data.

Goodhart's Law The observation that when a measure becomes an optimization target, it can stop representing the underlying goal.

GPT-2 A 1.5-billion parameter autoregressive language model released by OpenAI in 2019, serving as a primary lighthouse model for analyzing memory bandwidth constraints in transformer inference.

graceful degradation A system design principle where services continue functioning with reduced capabilities when faced with partial failures or data unavailability.

gradient accumulation A technique that simulates larger batch sizes by accumulating gradients from multiple smaller batches before updating model parameters, enabling training with limited memory.

gradient-based pruning A pruning method that uses gradient information to rank param-

eters, neurons, or filters as candidates for removal.

gradient clipping An optimization technique that mitigates exploding gradients by limiting their magnitude during backpropagation, typically by scaling gradients when their norm exceeds a threshold.

gradient compression A technique used in distributed training to reduce the communication overhead by compressing gradient information exchanged between computing nodes.

gradient descent An optimization algorithm that iteratively adjusts neural network parameters in the direction that minimizes the loss function, using gradients to determine update directions and magnitudes.

gradient synchronization The process in distributed training where locally computed gradients are aggregated across devices to ensure all devices update their parameters consistently.

GraNd (Gradient Normed) A data selection score based on the norm of an example's gradient during early training, identifying samples with larger learning signal.

graph break A point where a compiler cannot capture part of a program and must fall back to eager execution, splitting optimized graph regions.

graphics processing unit (GPU) A specialized processor originally designed for rendering graphics that provides massive parallel computing capabilities essential for efficient neural network computation, training, and inference.

Green AI A movement in AI research and practice that prioritizes computational efficiency and energy consumption as primary metrics alongside traditional performance metrics like accuracy.

Green500 Ranking system that evaluates the world's most powerful supercomputers based on energy efficiency measured in FLOP/s per watt rather than raw computational performance.

ground truth The reference target treated as correct for training or evaluation, often derived from measurement or annotation and therefore not necessarily identical to objective reality.

GRU Gated Recurrent Unit, a simplified variant of LSTM that uses fewer gates while maintaining the ability to capture long-term dependencies in sequential data.

H

hardware abstraction The layer in ML frameworks that provides a unified interface to diverse computing hardware (CPUs, GPUs, TPUs, accelerators) while handling device-specific optimizations and memory management behind the scenes.

hardware acceleration The use of specialized computing hardware to perform certain operations faster and more efficiently than software running on general-purpose processors.

hardware accelerator Specialized computing hardware designed to efficiently execute specific types of computations, such as GPUs for parallel processing or TPUs for machine learning workloads.

hardware-aware design The practice of designing neural network architectures specifically optimized for target hardware platforms, considering factors like memory hierarchy, compute units, and data movement patterns to maximize efficiency.

hidden layer An intermediate layer in a neural network between input and output layers that learns abstract representations by transforming data through learned weights and activation functions.

hidden state The internal memory of recurrent neural networks that carries information from previous time steps, enabling the network to maintain context across sequential inputs.

hierarchical processing A multi-tier system architecture where data and intelligence flow between different levels of the computing stack, from sensors to edge devices to cloud systems.

high-bandwidth memory (HBM) An advanced memory technology that provides much higher bandwidth than traditional DRAM by using 3D stacking and wide interfaces, critical for data-intensive AI workloads.

horizontal scaling Increasing system capacity by adding more machines or instances rather than upgrading existing hardware, providing better fault tolerance and load distribution.

hybrid machine learning The integration of multiple ML paradigms such as cloud, edge, mobile, and TinyML to form unified distributed systems that combine complementary strengths.

hybrid parallelism A distributed training approach that combines data parallelism and model parallelism to use the benefits of both strategies for training very large models.

hyperparameter A configuration setting that controls the learning process but is not learned from data, such as learning rate, batch size, or network architecture choices.

hyperscale data center Large-scale data center facilities containing thousands of servers and covering extensive floor space, designed to efficiently support massive computing workloads.

I

I/O bottleneck A performance limit where storage, decoding, preprocessing, or data transfer cannot supply data fast enough to keep compute resources busy.

idempotent transformation A data transformation that has the same effect when applied repeatedly as when applied once, supporting safe retries and reprocessing.

im2col A convolution-lowering technique that unfolds image patches into matrix columns so convolution can execute as GEMM.

ImageNet A massive visual database containing over 14 million labeled images across 20,000+ categories, developed by a team led by Fei-Fei Li and introduced in 2009, whose annual challenge became instrumental in driving breakthrough advances in computer vision.

imperative programming A programming paradigm in which programs specify sequences of commands that update state and direct control flow; it is distinct from whether operations execute eagerly or are staged for later execution.

inductive bias A structural assumption built into a model or architecture that restricts the functions it can learn efficiently, such as locality in CNNs.

inference The operational phase where trained neural networks make predictions on new data using fixed parameters, without weight updates.

InfiniBand A high-throughput, low-latency networking technology commonly used for

multi-node accelerator clusters and distributed training.

information-compute ratio (ICR) A metric quantifying the efficiency of data selection, defined as the ratio of model performance gain to the computational cost (FLOPs) required to achieve it. Maximizing ICR is the primary goal of data selection strategies.

infrastructure as code Practice of managing and provisioning computing infrastructure through machine-readable configuration files rather than manual processes, enabling version control and automation.

INT8 An 8-bit signed-integer representation used to quantize weights, activations, or both. Each INT8 value occupies one quarter the storage of an FP32 value, and supported hardware can accelerate applicable operations, with accuracy depending on the quantization scheme.

intermediate representation (IR) A compiler-internal representation between frontend model code and backend hardware code generation, used by ML frameworks for optimization and portability.

internet of things A network of physical objects embedded with sensors, software, and other technologies that connect and exchange data with other devices and systems over the internet.

intersectional analysis Evaluation that considers combinations of demographic attributes (for example, race and gender simultaneously) to detect concentrated harms not visible in single-factor analysis.

IOPS Input/Output Operations Per Second; a storage performance metric measuring the number of read/write operations a device can handle per second, critical for random access workloads like training data loading.

Iron Law of ML Systems A quantitative framework decomposing ML system performance into three terms: Data (limited by bandwidth), Compute (limited by FLOP/s), and Latency (limited by overhead). Formulated as $T = D_{vol}/BW + O/(R_{peak} \cdot \eta_{hw}) + L_{lat}$.

iterative pruning A gradual pruning strategy that removes parameters in multiple stages with fine-tuning between each stage, allowing the model to adapt to reduced capacity and typically achieving better accuracy than one-shot pruning.

J

JAX A numerical computing library developed by Google Research that combines NumPy's API with functional programming transformations including automatic differentiation, just-in-time compilation, and automatic vectorization for high-performance machine learning research.

JIT compilation Just-In-Time compilation that analyzes and optimizes code at runtime, enabling frameworks to balance the flexibility of eager execution with the performance benefits of graph optimization by compiling frequently used functions.

K

k-anonymity A privacy technique that ensures each record in a dataset is indistinguishable from at least k-1 other records by generalizing quasi-identifiers.

kernel A small matrix of learnable weights used in convolutional layers to detect specific features through the convolution operation, also called a filter. In GPU programming,

the term also denotes a function launched for parallel execution by a grid of threads.

kernel fusion An optimization technique that combines multiple computational operations into a single kernel to reduce memory transfers and improve performance on parallel processors.

keyword spotting (KWS) A technology that detects specific wake words or phrases in audio streams, typically used in voice-activated devices with constraints on power consumption and latency.

knowledge distillation A model compression technique where a smaller “student” network learns to mimic the behavior of a larger “teacher” network by training on the teacher’s soft output probabilities rather than just hard labels.

Kolmogorov-Smirnov test (K-S test) A non-parametric test that measures the maximum distance between two cumulative distributions, often used to compare feature distributions for drift.

Kullback-Leibler divergence (KL divergence) An asymmetric information-theoretic measure of how one probability distribution differs from another, used in monitoring and drift analysis.

KV cache Stored key and value tensors from previous transformer tokens used during autoregressive inference to avoid recomputing attention history.

L

L0 norm The count of nonzero elements in a vector or tensor, used in pruning to express sparsity or a target budget for remaining weights.

label quality drift Degradation in the reliability or consistency of labels over time, even when the underlying input distribution remains stable.

label shift A distribution shift where label frequencies change while the feature patterns conditioned on labels remain stable.

LAPACK Linear Algebra Package that extends BLAS with higher-level linear algebra operations including matrix decompositions, eigenvalue problems, and linear system solutions, providing essential mathematical foundations for machine learning computations.

latency The time delay between a request or event and the corresponding response or completion, critical in real-time applications where prompt responses are required.

latency constraints Real-time requirements that limit the maximum acceptable delay for model inference, driving optimization decisions in deployment scenarios where response time is critical.

layer normalization A normalization technique that normalizes inputs across the features dimension for each sample, commonly used in transformer architectures to stabilize training.

layerwise quantization A quantization granularity where all parameters within a single layer share the same quantization parameters, providing computational efficiency but potentially limiting representational precision compared to finer-grained approaches.

learning rate A hyperparameter that determines the step size for weight updates during gradient descent optimization, critically affecting training stability and convergence speed.

learning rate scheduling The systematic adjustment of learning rates during training, using strategies like step decay, exponential decay,

or cosine annealing to improve convergence and final model performance.

lifecycle coherence The principle that all stages of ML development should align with overall system objectives, maintaining consistency in data handling, model architecture, and evaluation criteria.

linear scaling failure The phenomenon where increasing computing resources (for example, doubling accelerators) results in sub-linear performance gains due to communication overhead and synchronization costs.

LINPACK Benchmark developed at Argonne National Laboratory that measures system performance by solving dense systems of linear equations, famous for its use in TOP500 supercomputer rankings.

Little’s Law A queuing relationship stating that average concurrency equals arrival rate times average time in the system, linking throughput, latency, and capacity.

Llama Large Language Model Meta AI, a family of open foundation models that popularized efficient architectural choices like RMSNorm and SwiGLU, serving as a modern reference point for large-scale transformer efficiency.

load balancing Distribution of incoming requests across multiple server instances to prevent bottlenecks, improve response times, and ensure high availability.

locality-sensitive hashing (LSH) A hashing family that maps similar items to the same buckets with high probability, enabling scalable approximate similarity search.

logits The raw, unnormalized scores output by the last layer of a neural network before the activation function (like softmax) converts them into probabilities.

loss function A mathematical function that measures how well model outputs satisfy a training objective, providing the objective minimized by the training algorithm.

loss scaling A technique used in mixed-precision training that multiplies the loss by a large factor before backpropagation to prevent gradient underflow in reduced precision formats.

Lottery Ticket Hypothesis The theory that large neural networks contain sparse subnetworks that, when trained in isolation from proper initialization, can achieve comparable accuracy to the full network while being significantly smaller.

low-rank factorization A matrix decomposition technique that approximates large weight matrices as products of smaller matrices, reducing the number of parameters and computational operations required for neural network layers.

LSTM Long Short-Term Memory, a type of recurrent neural network architecture designed to handle long-term dependencies through gating mechanisms that control information flow.

M

machine learning A subset of artificial intelligence that enables systems to automatically improve performance on tasks through experience and data rather than explicit programming.

machine learning framework A software platform that provides tools and abstractions for designing, training, and deploying machine learning models, bridging user applications with infrastructure through computational graphs, hardware optimization, and workflow orchestration. Examples include TensorFlow, PyTorch, and JAX.

machine learning lifecycle A structured, iterative process that encompasses all stages involved in developing, deploying, and maintaining machine learning systems, from problem definition through ongoing monitoring and improvement.

machine learning operations (MLOps) The engineering discipline and set of tools focused on operationalizing machine learning models through automation, monitoring, and management of the entire ML pipeline from development to production.

machine learning systems engineering The engineering discipline focused on building reliable, efficient, and scalable AI systems across computational platforms, spanning the entire AI lifecycle from data acquisition through deployment and operations with emphasis on resource-awareness and system-level optimization.

macro benchmarks Evaluation methodology that assesses complete machine learning models to understand how architectural choices and component interactions affect overall system behavior and performance.

magnitude-based pruning The most common pruning method that removes parameters with the smallest absolute values, based on the assumption that weights with smaller magnitudes contribute less to the model’s output.

masking An anonymization technique that alters or obfuscates sensitive values so they cannot be directly traced back to the original data subject.

membership inference attack A privacy attack that tries to determine whether a specific record was included in a model’s training data.

memory bandwidth Rate at which data can be read from or written to memory, measured in bytes per second, which often becomes a bottleneck in memory-intensive machine learning workloads.

memory hierarchy The organization of memory systems with different access speeds and capacities, from fast on-chip caches to slower off-chip main memory.

memory wall The widening performance gap between processor speed and memory bandwidth, where computational capacity outpaces the rate at which data can be delivered to the processor, becoming a primary bottleneck for large ML models.

metadata Descriptive information about datasets that includes details about data collection, quality metrics, validation status, and other contextual information essential for data management.

micro-batch A smaller slice of an effective batch processed sequentially for gradient accumulation or pipeline parallelism to reduce peak memory.

micro benchmarks Specialized evaluation tools that assess individual components or specific operations within machine learning systems, such as tensor operations or neural network layers.

microcontroller A small computer on a single integrated circuit containing a processor core, memory, and programmable input/output peripherals, commonly used in embedded systems.

MinHash A sketching method that approximates set similarity with compact signatures, commonly used for scalable near-duplicate detection.

mini-batch gradient descent A training approach that computes gradients and updates weights using a small subset of training examples simultaneously, balancing compu-

- tational efficiency with gradient estimation quality.
- mini-batch processing** An optimization approach that computes gradients over small batches of examples, balancing the computational efficiency of batch processing with the memory constraints of stochastic methods.
- mixed-precision computing** A technique that uses different numerical precisions at various stages of computation, such as FP16 for matrix multiplications and FP32 for accumulations.
- mixed-precision training** A training methodology that combines different numerical precisions (typically FP16 and FP32) to optimize memory usage and computational speed while maintaining training stability.
- ML node** A single production ML application treated as an operational unit, including data pipelines, feature computation, model training, serving infrastructure, and monitoring.
- ML systems** Integrated computing systems comprising three core components: data that guides algorithmic behavior, learning algorithms that extract patterns from data, and computing infrastructure that enables both training and inference processes.
- ML systems spectrum** The range of machine learning system deployments from cloud-based systems with abundant resources to tiny embedded devices with severe constraints, each requiring different optimization strategies and trade-offs.
- ML Test Score** A production-readiness rubric for ML systems that evaluates data, model, infrastructure, and monitoring tests.
- MLCommons** Organization that develops and maintains industry-standard benchmarks for machine learning systems, including the MLPerf suite for training and inference evaluation.
- MLPerf** Industry-standard benchmark suite that provides standardized tests for training and inference across various deep learning workloads, enabling fair comparisons of machine learning systems.
- MLPerf execution scenarios** Standard MLPerf inference traffic scenarios that model sequential, synchronized, server, batch, and interactive workloads with different latency and throughput requirements.
- MLPerf Inference** Benchmark framework that evaluates machine learning inference performance across different deployment environments, from cloud data centers to mobile devices and embedded systems.
- MLPerf Mobile** Specialized benchmark that extends MLPerf evaluation to smartphones and mobile devices, measuring latency and responsiveness under strict power and memory constraints.
- MLPerf Power** An MLPerf methodology for measuring power and energy efficiency of ML workloads alongside performance.
- MLPerf Tiny** Benchmark designed for embedded and ultra-low-power AI systems such as IoT devices, wearables, and microcontrollers operating with minimal processing capabilities.
- MLPerf Training** Standardized benchmark that evaluates machine learning training performance by measuring time-to-accuracy, throughput, and resource utilization across different hardware platforms.
- mobile machine learning** The execution of machine learning models directly on portable, battery-powered devices like smartphones and tablets, enabling personalized and responsive applications.
- model cards** A standardized format for documenting machine learning models, capturing information essential for responsible deployment, including intended use, performance factors, and ethical considerations.
- model-centric AI** A research paradigm where the dataset is treated as fixed and engineering effort focuses on optimizing model architecture.
- model compression** Techniques used to reduce the size and computational requirements of machine learning models while seeking to preserve accuracy, enabling deployment on resource-constrained devices.
- model deployment** The process of integrating trained machine learning models into production systems where they can make predictions on new data and provide value to end users.
- model drift** The degradation of machine learning model performance over time due to changes in data patterns, user behavior, or environmental conditions that differ from the original training conditions.
- model evaluation** The systematic assessment of machine learning model performance using various metrics and validation techniques to determine whether the model meets requirements and is ready for deployment.
- model FLOPs utilization (MFU)** The ratio of the model FLOPs per second implied by observed throughput to the hardware's theoretical peak FLOPs per second, measuring how effectively a workload uses available compute for model computation.
- model optimization** The systematic refinement of machine learning models to enhance their efficiency while maintaining effectiveness, balancing trade-offs between accuracy, computational cost, memory usage, latency, and energy efficiency.
- model parallelism** A distributed training strategy that splits a neural network model across multiple devices, with each device responsible for computing a portion of the network.
- model quantization** The process of reducing numerical precision in machine learning models, typically from FP32 to INT8, to decrease model size and memory traffic; speed increases only with efficient low-precision kernels.
- model registry** Centralized repository for storing, versioning, and managing trained machine learning models with associated metadata, facilitating model governance and deployment.
- model serving** Infrastructure and systems that expose deployed machine learning models through APIs to handle prediction requests at scale with appropriate latency and throughput.
- model training** The process of using machine learning algorithms to learn patterns from training data, adjusting model parameters to minimize prediction errors and create a functional predictive system.
- model validation** The process of testing machine learning models on independent datasets to assess their generalization ability and ensure they perform reliably on unseen data.
- model versioning** The systematic tracking and management of different versions of machine learning models, including their parameters, training data, and performance metrics, to enable comparison and rollback capabilities.
- momentum** An optimization technique that accumulates a velocity vector across iterations to help gradient descent navigate through local minima and accelerate convergence in consistent gradient directions.
- monitoring** The continuous observation and measurement of machine learning system performance, data quality, and operational metrics in production to detect issues and trigger maintenance actions.
- multi-head attention** An attention mechanism that uses multiple parallel attention heads, each focusing on different aspects of the input to capture diverse types of relationships simultaneously.
- multilayer perceptron (MLP)** A feedforward neural network with one or more hidden layers between input and output layers, capable of learning nonlinear mappings through dense connections and activation functions.
- multiply-accumulate (MAC)** The operation of multiplying two values and accumulating the result into a running sum, dominating linear layers and accelerator datapaths.
- MYCIN** One of the first large-scale expert systems developed at Stanford beginning in 1972 to diagnose blood infections, representing the shift toward capturing human expert knowledge in specific domains rather than pursuing general artificial intelligence.

N

- N:M sparsity** A structured sparsity pattern where exactly N of every M values are nonzero, giving hardware more regular zero patterns to exploit.
- neural architecture search** An automated approach that searches possible combinations of layers, connections, and architectural hyperparameters to find neural network architectures that perform well under specified objectives and constraints.
- neural network** A computational model consisting of interconnected nodes organized in layers that can learn to map inputs to outputs through adjustable connection weights.
- neural processing unit (NPU)** Specialized processors designed specifically for accelerating neural network operations and machine learning computations, optimized for parallel processing of AI workloads.
- neuromorphic computing** A computing approach that mimics the structure and function of biological neural networks, potentially offering more energy-efficient processing for AI applications.
- NoSQL** A category of database systems designed to handle large volumes of unstructured or semi-structured data with flexible schemas, often used in big data applications.
- numerical precision optimization** The dimension of model optimization that addresses how numerical values are represented and processed, including quantization techniques that map high-precision values to lower-bit representations.
- NVLink** NVIDIA's high-bandwidth GPU interconnect for intra-node communication, useful for gradient synchronization and model-parallel activation transfers.
- observability** Comprehensive monitoring approach that provides insight into system behavior through metrics, logs, and traces, enabling understanding of internal states from external outputs.
- on-chip memory** Fast memory integrated directly onto the processor chip, including caches and scratchpad memory, providing high bandwidth and low latency data access.
- on-device learning** The capability for machine learning models to adapt and learn directly on edge devices without requiring data transmission to external servers.

one-hot encoding A representation where categorical labels become vectors with a single 1 and remaining 0s.

one-shot pruning A pruning strategy where a large fraction of parameters is removed in a single step, typically followed by fine-tuning to recover accuracy, offering simplicity but potentially requiring more aggressive fine-tuning.

online analytical processing (OLAP) A database approach optimized for complex analytical queries across large datasets, typically used in data warehouses for business intelligence.

online inference Real-time prediction serving that processes individual requests with low latency, suitable for interactive applications requiring immediate responses.

online transaction processing (OLTP) A database approach optimized for frequent, short transactions and real-time processing, commonly used in operational applications.

ONNX Open Neural Network Exchange, a standardized format for representing machine learning models that enables interoperability between different frameworks, allowing models trained in one framework to be deployed using another.

ONNX Runtime Cross-platform inference engine that optimizes machine learning models through techniques like operator fusion and kernel tuning to improve inference speed and reduce computational overhead.

optimizer An algorithm that adjusts model parameters during training to minimize the loss function, with common examples including SGD (Stochastic Gradient Descent), Adam, and RMSprop, each with different strategies for parameter updates.

optimizer state Auxiliary values maintained by an optimizer in addition to model parameters, such as Adam's momentum and variance tensors, often dominating training memory for large models.

orchestration Coordination and management of complex workflows and distributed computing tasks, often using platforms like Kubernetes for container management.

outlier detection The process of identifying data points that significantly deviate from normal patterns, which may represent errors, anomalies, or valuable rare events.

overfitting A phenomenon where a model learns specific details of training data so well that it fails to generalize to new, unseen examples, typically indicated by high training accuracy but poor validation performance.

P

padding A technique in convolutional networks that adds zeros or other values around the input borders to control the spatial dimensions of the output feature maps.

PagedAttention A KV-cache memory-management technique that stores attention context in fixed-size pages rather than contiguous per-request blocks to reduce fragmentation.

paradigm shift A fundamental change in scientific approach, like the shift from symbolic reasoning to statistical learning in AI during the 1990s, and from shallow to deep learning in the 2010s, requiring researchers to abandon established methods for radically different approaches.

parallelism The simultaneous execution of multiple computational tasks or operations, fundamental to achieving high performance in neural network processing.

parameter A learnable component of a neural network, including weights and biases, that gets adjusted during training to minimize the loss function.

Pareto frontier The set of configurations where improving one objective requires worsening another, used to reason about accuracy, latency, memory, and energy trade-offs.

partitioning A database technique that divides large datasets into smaller, manageable segments based on specific criteria to improve query performance and system scalability.

perceptron The fundamental building block of neural networks, consisting of weighted inputs, a bias term, and an activation function that produces a single output.

performance insights Analytical observations derived from monitoring production machine learning systems that reveal opportunities for improvement in model accuracy, system efficiency, or user experience.

performance-per-data (PPD) A metric measuring the accuracy gain per training sample, used to evaluate the quality of a dataset or selection strategy. High PPD indicates a dataset rich in information and low in redundancy.

pinned memory Page-locked host memory that enables direct, asynchronous transfers between CPU and GPU, improving data-loading throughput at the cost of reserving system memory.

pipeline jungle Anti-pattern where complex, interdependent data processing pipelines become difficult to maintain, debug, and modify, leading to technical debt and operational complexity.

pipeline parallelism A form of model parallelism where different layers of a model are placed on different devices and data flows through them in a pipeline fashion, allowing multiple batches to be processed simultaneously.

point-in-time correctness The guarantee that training examples use only feature values that would have been available at the historical prediction time, preventing leakage.

pooling A downsampling operation in convolutional networks that reduces spatial dimensions while retaining important features, commonly using max or average operations over local regions.

population stability index (PSI) A binned distribution-shift metric that compares current feature distributions against a baseline to detect population drift.

positional encoding A method used in transformer architectures to inject information about the position of tokens in a sequence, since transformers lack inherent sequential processing.

post-training quantization A quantization approach applied to already-trained models without modifying the training process, typically involving calibration on representative data to determine optimal quantization parameters.

power usage effectiveness (PUE) Metric used in data centers to measure energy efficiency, calculated as the ratio of total facility power consumption to IT equipment power consumption.

power wall The technological barrier reached around 2005 where increasing processor frequency no longer yielded performance gains without unsustainable increases in power density and heat generation, forcing a shift to parallel and specialized architectures.

precision In numerical computing, the number of significant bits or digits used to represent values, affecting computational accu-

racy and resource requirements in machine learning systems.

prefetching A system optimization technique that loads data into memory before it is needed, overlapping data loading with computation to reduce idle time and improve training throughput.

problem definition The initial stage of machine learning development that involves clearly specifying objectives, constraints, success metrics, and operational requirements to guide all subsequent development decisions.

processing element (PE) A repeated accelerator building block containing compute units and local storage, with arrays of PEs providing parallel tensor execution.

programmable logic controller (PLC) Industrial control systems used in manufacturing and IoT environments that can be integrated with ML models for automated decision-making in operational technology contexts.

protein folding problem The scientific challenge of predicting the three-dimensional structure and dynamics of proteins from their amino acid sequences and environments, a longstanding problem on which systems such as AlphaFold achieved breakthrough accuracy without resolving every folding question.

proxy variable A feature that indirectly carries signal about a protected or sensitive attribute even when that attribute is removed.

pseudonymization A privacy technique that replaces direct identifiers with artificial identifiers while maintaining the ability to trace records for analysis purposes.

PyTorch A deep learning framework developed by Facebook's AI Research lab that emphasizes dynamic computational graphs, eager execution, and intuitive Python integration, particularly popular for research and experimentation.

Q

quantization A model-compression technique that represents parameters or activations at lower precision, such as FP32 to INT8, reducing storage and data movement without reducing mathematical operation count; speedups require supported hardware and kernels.

quantization-aware training A training approach where quantization effects are simulated during the training process, allowing the model to adapt to reduced precision and typically achieving better accuracy than post-training quantization.

quantization granularity The level at which quantization parameters are applied, ranging from per-tensor (coarsest) to per-channel or per-group (finer), with finer granularity typically preserving more accuracy but requiring more storage.

queries per second (QPS) Performance metric that measures how many inference requests a system can process in one second, commonly used to evaluate throughput in production deployments.

query-key-value The three components of attention mechanisms where queries determine what to look for, keys represent what is available, and values contain the actual information to be weighted and combined.

R

real-time processing Processing data as it becomes available under application timing

constraints; hard real-time systems guarantee worst-case deadlines, while soft real-time systems tolerate occasional misses.

receptive field The region of the input that influences a particular neuron's output, determining the spatial extent of patterns that can be detected by that neuron.

recurrent neural network A type of neural network designed for sequential data processing, featuring connections that create loops allowing information to persist across time steps.

Red AI AI research and development that prioritizes maximizing accuracy or performance without regard for the increasing computational and environmental costs required.

regularization Techniques used to prevent overfitting in neural networks by adding constraints or penalties, including methods like dropout, weight decay, and data augmentation.

ReLU Rectified Linear Unit activation function defined as $f(x) = \max(0, x)$ that introduces nonlinearity while maintaining computational efficiency and helping mitigate vanishing gradients for positive inputs.

residual connection A skip connection that adds the input of a layer to its output, enabling the training of very deep networks by mitigating the vanishing gradient problem.

ResNet Residual Network, a deep convolutional architecture built from residual blocks with skip connections, enabling the training of networks with hundreds of layers and achieving breakthrough performance.

ResNet-50 A specific 50-layer variant of the Residual Network architecture that serves as the canonical lighthouse model for compute-bound workloads, balancing depth and computational cost for vision benchmarks.

retinal fundus photographs Medical images of the interior surface of the eye, including the retina, optic disc, and blood vessels, commonly used for diagnosing eye diseases and training medical AI systems.

reverse-mode differentiation An automatic differentiation technique that computes gradients by traversing the computational graph in reverse order, highly efficient for functions with many inputs and few outputs, making it ideal for neural network training.

ridge point The arithmetic intensity at which a workload transitions from memory-bound to compute bound on a given hardware platform, calculated as peak FLOP/s divided by memory bandwidth. The inflection point on the Roofline Model diagram.

Ring AllReduce A bandwidth-efficient AllReduce algorithm where devices pass chunks around a ring and accumulate results for distributed gradient synchronization.

RMSprop An adaptive learning rate optimization algorithm that maintains a moving average of squared gradients to automatically adjust learning rates for each parameter during training.

role-based access control (RBAC) An access-control model that grants permissions according to user, service, or team roles rather than ad hoc individual grants.

rollback Process of reverting to a previous stable version of a model or system when issues are detected in production, ensuring service continuity.

roofline analysis A performance modeling technique that plots attainable performance against operational intensity, with peak compute and bandwidth forming the ceilings, to identify whether a system is memory bound or compute bound.

S

scalability The ability of machine learning systems to handle increasing amounts of data, users, or computational demands without significant degradation in performance or user experience.

scale and zero-point Quantization parameters that map real-valued tensors to integer values and back, where scale controls step size and zero-point maps real zero into the quantized range.

schema The structure and format definition of data that specifies data types, field names, and relationships, essential for data validation and processing consistency.

schema evolution The process of modifying data schemas over time while maintaining backward compatibility and ensuring continued functionality of dependent systems and applications.

schema-on-read An approach used in data lakes where data structure is defined and enforced at the time of reading rather than when storing, providing flexibility for diverse data types.

scratchpad memory Fast software-managed local memory used by accelerators to stage predictable data movement under compiler or runtime control.

segmentation maps Detailed annotations that classify objects at the pixel level, providing the most granular labeling information but requiring significantly more storage and processing resources.

self-attention An attention mechanism where queries, keys, and values all come from the same sequence, allowing each position to attend to all positions including itself.

semi-supervised learning A machine learning approach that uses both labeled and unlabeled data for training, using structural assumptions to improve model performance with limited labels.

sensitivity The proportion of actual positive cases correctly identified by a classification model, also known as true positive rate or recall; critical in medical applications where missing positive cases has severe consequences.

serverless Cloud computing model where infrastructure is automatically managed by the provider, allowing code execution without server management concerns.

service level agreement (SLA) A formal external contract specifying service commitments, often through one or more SLOs, and the consequences or penalties for noncompliance.

service-level objective (SLO) An internal measurable target for a service-level indicator such as latency, error rate, or availability that guides operations and may underpin an external SLA.

shadow deployment A deployment strategy where a new model runs in parallel with the production model, making predictions that are logged but not served to users, enabling validation without user impact.

shallow learning Machine learning approaches that use algorithms with limited complexity, such as support vector machines and decision trees, which require carefully engineered features but cannot automatically discover hierarchical representations like deep learning methods.

SHAP values Feature-attribution scores based on Shapley values that estimate each feature's contribution to a model prediction.

sigmoid An activation function that maps input values to a range between 0 and 1, histori-

cally popular but prone to vanishing gradient problems in deep networks.

silent bias Model unfairness that produces valid-looking but discriminatory outputs, evading traditional error monitoring and requiring disaggregated evaluation to detect.

silent degradation The gradual decline in ML system performance that occurs without triggering errors, exceptions, or alerts. An ML system can continue operating while producing increasingly inaccurate predictions.

silent failure A system failure mode where an ML model continues to produce plausible-looking outputs that are gradually less accurate or contextually relevant without triggering conventional error alerts.

SIMD (Single Instruction, Multiple Data) A parallel computing architecture that applies the same operation to multiple data elements simultaneously, effective for regular data-parallel computations.

SIMT (Single Instruction, Multiple Threads) An extension of SIMD that enables parallel execution across multiple independent threads, each maintaining its own state and program counter.

singular value decomposition A matrix factorization technique that decomposes a matrix into the product of three matrices, commonly used in low-rank approximations to compress neural network layers by retaining only the most significant singular values.

skip connection A direct connection that bypasses one or more layers, allowing gradients to flow more easily through deep networks and enabling better training of very deep architectures.

soft targets Teacher-model probability distributions used as supervision in knowledge distillation, carrying class-similarity information beyond hard labels.

softmax An activation function that converts logits into a probability distribution where outputs sum to 1, used in multi-class classification.

sparse Tensor Core A Tensor Core variant that accelerates supported structured sparsity patterns such as 2:4 sparsity by skipping constrained zero values.

sparsity The property of tensors in which many values are exactly zero, which can improve computational efficiency when the zeros have suitable structure and specialized sparse kernels or hardware support them; near-zero values count only after pruning.

SPEC CPU Standardized benchmark suite developed by the System Performance Evaluation Cooperative that measures processor performance using real-world applications rather than synthetic tests.

special function unit (SFU) Dedicated hardware for specialized operations such as exponentials, reciprocals, roots, and some activation-function primitives that do not map cleanly to ordinary arithmetic or matrix multiply units.

specificity The proportion of actual negative cases correctly identified by a classification model, also known as true negative rate; important for avoiding overwhelming referral systems with false positives.

SPECpower Benchmark methodology that measures server energy efficiency across varying workload levels, enabling direct comparisons of power-performance trade-offs in computing systems.

speculative decoding An optimization technique for autoregressive language models where a smaller model generates draft tokens that are then verified by a larger model,

- accelerating inference while maintaining quality.
- speed of light (latency)** The physical speed limit of information transmission (approx. 200,000 km/s in fiber), creating an irreducible lower bound on network latency that necessitates edge computing for applications requiring sub-10 ms response times over long distances.
- state dict** A framework serialization mapping from module or optimizer names to tensors, used to save, load, and checkpoint models and optimizer state.
- static graph** A computational graph that is defined completely before execution begins, enabling comprehensive optimization and efficient deployment but requiring all operations to be specified upfront, limiting runtime flexibility.
- static quantization** A quantization approach where quantization parameters are determined once during calibration and remain fixed during inference, providing computational efficiency but less adaptability than dynamic approaches.
- statistical learning** The era of machine learning that emerged in the 1990s, shifting focus from rule-based symbolic AI to algorithms that could learn patterns from data, laying the groundwork for modern data-driven approaches to artificial intelligence.
- stochastic gradient descent** A variant of gradient descent that estimates gradients using individual training examples or small batches rather than the entire dataset, reducing memory requirements and enabling online learning.
- Straight-Through Estimator (STE)** A gradient approximation that lets training pass gradients through nondifferentiable operations such as rounding, often used in quantization-aware training.
- stream ingestion** A data processing pattern that handles data in real-time as it arrives, essential for applications requiring immediate processing and low-latency responses.
- stream processing** Real-time data processing approach that handles continuous flows of data as it arrives, enabling immediate responses to events and pattern detection.
- stride** The step size by which a convolutional filter moves across the input, controlling the spatial dimensions of the output and the degree of overlap between filter applications.
- structured pruning** A pruning approach that removes entire computational units such as neurons, channels, or layers, producing smaller dense models that are more hardware-friendly than the sparse matrices created by unstructured pruning.
- supervised learning** A machine learning approach where models learn from labeled training examples to make predictions on new, unlabeled data.
- symbolic AI** An approach to artificial intelligence that uses high-level symbolic representations of problems, logic, and search algorithms, dominant before the deep learning revolution and characterized by expert systems and rule-based reasoning.
- synthetic benchmark** Artificial test program designed to measure specific aspects of system performance, as opposed to benchmarks based on real-world applications and workloads.
- synthetic data** Artificially generated data created using algorithms, simulations, or generative models to supplement real-world datasets, addressing limitations in data availability or privacy concerns.
- system entropy** The tendency of ML systems to degrade over time as the world changes (drift) or as hidden dependencies accumulate (technical debt), requiring active energy (ops) to maintain order.
- System-on-Chip (SoC)** An integrated circuit that incorporates most or all components of a computer or electronic system, including CPU, GPU, memory, and specialized processors on a single chip. Commonly used in mobile devices and embedded systems for space and power efficiency.
- systems co-design** Joint design of algorithms, software, and hardware so model structure and execution fit physical capabilities and constraints.
- systems integration** The process of combining various components and subsystems into a unified, functional system that operates efficiently and reliably as a whole.
- systems thinking** An approach to understanding complex systems by considering how individual components interact and affect the whole system, particularly important in ML where data, algorithms, hardware, and deployment environments must work together effectively.
- systolic array** A specialized hardware architecture that efficiently performs matrix operations by streaming data through a grid of processing elements, minimizing data movement and energy consumption.
- T**
- tail latency** High-percentile response time in a system, typically measured as P95 or P99 latency, important for characterizing slow requests under peak load but not equivalent to the absolute worst case. It complements median and mean latency rather than replacing them.
- tailored inference benchmarks** Specialized performance tests designed for specific deployment environments or use cases, accounting for unique constraints and optimization requirements.
- tanh** Hyperbolic tangent activation function that maps inputs to (-1, 1), providing zero-centered outputs.
- technical debt** Long-term maintenance cost accumulated from expedient design decisions during development, particularly problematic in ML systems due to data dependencies and model complexity.
- telemetry** Automated collection and transmission of performance data and metrics from distributed systems, enabling remote monitoring and analysis.
- tensor** A multi-dimensional array used to represent data in neural networks, generalizing scalars (0D), vectors (1D), and matrices (2D) to higher dimensions.
- Tensor Core** A specialized GPU matrix unit for accelerated mixed-precision multiply-accumulate operations on tensor tiles.
- tensor decomposition** The extension of matrix factorization to higher-order tensors, used to compress neural network layers by representing weight tensors as combinations of smaller tensors with fewer parameters.
- tensor parallelism** A model parallelism strategy where individual tensor operations (like matrix multiplication) are split across multiple devices, reducing memory per device and potentially reducing large-layer latency when the reduction in computation time exceeds the added communication time.
- tensor processing unit (TPU)** Google's custom application-specific integrated circuit designed specifically for machine learning workloads, optimized for matrix operations and featuring systolic array architecture.
- TensorFlow** A comprehensive machine learning framework developed by Google that provides tools for the entire ML pipeline from research to production, featuring both eager execution and graph-based computation with extensive ecosystem support.
- TensorRT** NVIDIA's inference optimization library that applies techniques like operator fusion and precision reduction to accelerate deep learning inference on GPU hardware.
- ternarization** An extreme quantization technique that constrains weights to three values (typically -1, 0, +1), providing significant compression while maintaining more representational capacity than binary quantization.
- thermal design power (TDP)** The sustained power and heat level a device or cooling system is designed to dissipate under typical maximum workload.
- thermal throttling** A protective mechanism in mobile and embedded devices that reduces processor clock speed and performance to prevent overheating, often limiting the sustained performance of on-device ML inference.
- throughput** The rate at which a system can process data or complete operations, typically measured in operations per second and crucial for training large models.
- time per output token (TPOT)** The per-token latency during autoregressive decode after the first token, reflecting decode throughput and memory-bandwidth limits.
- time-to-accuracy** The wall-clock time required to train a model to a specified validation accuracy.
- time to first token (TTFT)** The time from request arrival to the first generated token in streaming LLM serving, including prefill, scheduling, and initial response latency.
- TinyML** A field focused on deploying machine learning models on ultra-constrained embedded devices such as microcontrollers and sensors, operating in the milliwatt to sub-watt power range, with severe limitations on memory, power, and computational capacity. Also called *tiny machine learning*.
- TOPS** Tera Operations Per Second, an operation-rate unit indicating how many trillion operations a system can execute in one second. It is commonly reported for low-precision integer workloads, so the operation type and precision must be specified.
- total cost of ownership (TCO)** A comprehensive financial metric for ML systems encompassing training, inference, and operational costs over the system's entire lifecycle.
- train-serve split** A hybrid architecture pattern where computationally intensive model training occurs on powerful cloud infrastructure, while the trained model is optimized and deployed for inference on resource-constrained edge or mobile devices.
- training** The process of adjusting neural network parameters using data or interaction experience and an optimization procedure to minimize a loss or maximize a reward.
- training-serving skew** A mismatch between how features or data are computed during model training vs. serving in production, which can degrade model performance. Common causes include different preprocessing pipelines, feature computation timing, or data sources between training and inference.
- transfer learning** A machine learning technique that uses knowledge gained from pretrained models on related tasks, allowing faster training and better performance on new

tasks with limited data by reusing learned features and representations.

transformer A neural network architecture built from attention, feed-forward blocks, residual connections, and normalization; the original design eliminates recurrence and convolution, though later variants may incorporate them.

translation equivariance A property where shifting the input causes a corresponding shift in the output feature map, distinct from invariance which discards position.

translation invariance The property that a model's output remains unchanged when the input shifts. Convolution is translation-equivariant; pooling or global aggregation can produce approximate invariance.

Tucker decomposition A tensor decomposition method that generalizes singular value decomposition to higher-order tensors using a core tensor and factor matrices, commonly used for compressing convolutional neural network layers.

TV white spaces Unused broadcasting frequencies that can be repurposed for internet connectivity, as employed by systems like FarmBeats to extend network access to remote agricultural sensors and IoT devices.

U

uniform quantization A quantization approach where the range of values is divided into evenly spaced intervals, providing simple implementation but potentially suboptimal for nonuniform value distributions.

universal approximation theorem A theoretical result proving that neural networks with sufficient width and a suitable nonlinear activation function can approximate any continuous function on a compact domain to arbitrary accuracy.

unreasonable effectiveness of data The empirical observation that for many problems, adding more data is more effective than improving algorithms, driving the data-centric AI paradigm.

unstructured pruning A pruning approach that removes individual weights while preserving the overall network architecture, creating sparse weight matrices that require specialized hardware support to realize computational benefits.

unstructured sparsity A form of model sparsity where individual weights are set to zero without following any particular pattern, creating irregular sparsity patterns that require specialized hardware support to realize computational benefits.

V

validation issues Problems identified during model testing that indicate poor performance, overfitting, data quality problems, or other issues that must be resolved before deployment.

vanishing gradient problem A problem in deep neural networks where gradients become exponentially smaller as they propagate backward through layers, making it difficult for early layers to learn effectively.

vector operations Computational operations that process multiple data elements simultaneously, enabling efficient parallel execution of element-wise transformations in neural networks.

verification gap The mismatch between finite test-set coverage and the much larger input space an ML system may encounter in production.

versioning The practice of tracking changes to datasets, models, and pipelines over time, enabling reproducibility, rollback capabilities, and audit trails in ML systems.

virtuous cycle The self-reinforcing process in deep learning where improvements in data availability, algorithms, and computing power each enable further advances in the other areas, accelerating overall progress.

Von Neumann bottleneck The performance limitation caused by the shared bus between processor and memory in traditional computer architectures, where data movement becomes more expensive than computation.

W

warehouse-scale computer (WSC) A data center operated as a single programmable computer, treating compute, storage, networking, and failures as system-level resources.

warp A group of 32 threads on NVIDIA GPUs whose active threads execute one common instruction at a time; the fundamental unit of thread scheduling on NVIDIA GPUs.

Waymo A subsidiary of Alphabet Inc. that represents a leading deployment of machine learning systems in autonomous vehicle technology, demonstrating how ML systems can span from embedded systems to cloud infrastructure in safety-critical environments.

weak supervision An approach that uses lower-quality labels obtained more efficiently through heuristics, distant supervision, or programmatic methods rather than manual expert annotation.

web scraping An automated technique for extracting data from websites to build custom datasets, requiring careful consideration of legal, ethical, and technical constraints.

weight A learnable parameter that determines the strength of connection between neurons in different layers, adjusted during training to minimize the loss function.

weight matrix An organized collection of weights connecting one layer to another in a neural network, enabling efficient computation through matrix operations.

weight sharing The practice of using the same parameters across different spatial locations, as in convolutional networks, reducing the number of parameters while maintaining pattern detection capabilities.

Whetstone Early benchmark developed at the UK National Physical Laboratory by Curnow and Wichmann, first published as an ALGOL 60 program in 1972 and later as the canonical FORTRAN version in 1976. It measured floating-point arithmetic performance in MWIPS (millions of Whetstone instructions per second), becoming one of the first widely-adopted standardized performance tests.

workflow orchestration Automated coordination and management of complex ML pipeline sequences, ensuring proper execution order, dependency management, and error handling across distributed systems.

X

XLA Accelerated Linear Algebra, a domain-specific compiler for linear algebra operations that optimizes TensorFlow and JAX computations by generating efficient code for various hardware platforms including CPUs, GPUs, and TPUs.

References

- Abadi, Martin, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, et al. 2016. "TensorFlow: A System for Large-Scale Machine Learning." *Proceedings of the 12th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 265–83.
- Agarwal, Pankaj K., Sariel Har-Peled, and Kasturi R. Varadarajan. 2005. "Geometric Approximation via Coresets." In *Combinatorial and Computational Geometry*, edited by Jacob E. Goodman, János Pach, and Emo Welzl, vol. 52. MSRI Publications. Cambridge University Press. <https://doi.org/10.1017/9781009701259.002>.
- Agrawal, Amey, Nitin Kedia, Ashish Panwar, Jayashree Mohan, Nipun Kwatra, Bhargav S. Gulavani, Alexey Tumanov, and Ramachandran Ramjee. 2024. "Taming Throughput-Latency Tradeoff in LLM Inference with Sarathi-Serve." *18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24)*, 117–34.
- Ainslie, Joshua, James Lee-Thorp, Michiel de Jong, Yury Zemlyanskiy, Federico Lebron, and Sumit Sanghai. 2023. "GQA: Training Generalized Multi-Query Transformer Models from Multi-Head Checkpoints." *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 4895–901. <https://doi.org/10.18653/v1/2023.emnlp-main.298>.
- Akida, Tyler, Robert Bradshaw, Craig Chambers, Slava Chernyak, Rafael J. Fernández-Moctezuma, Reuven Lax, Sam McVeety, et al. 2015. "The Dataflow Model: A Practical Approach to Balancing Correctness, Latency, and Cost in Massive-Scale, Unbounded, Out-of-Order Data Processing." *Proceedings of the VLDB Endowment* 8 (12): 1792–803. <https://doi.org/10.14778/2824032.2824076>.
- Altini, Marco, and Hannu Kinnunen. 2021. "The Promise of Sleep: A Multi-Sensor Approach for Accurate Sleep Stage Detection Using the Oura Ring." *Sensors* 21 (13): 4302. <https://doi.org/10.3390/s21134302>.
- Amazon Web Services. 2024a. *Amazon Mechanical Turk*.
- Amazon Web Services. 2024b. *Amazon Simple Storage Service (S3)*.
- Amazon Web Services. 2024c. *Amazon Web Services (AWS)*.
- Amazon Web Services. 2024d. *AWS Auto Scaling*.
- Amazon Web Services. 2024e. *AWS CloudFormation*.
- Amazon Web Services. 2026. *Amazon EC2 Spot Instances*.
- AMD. 2023. *AMD Instinct MI300X Accelerators*. AMD product documentation.
- Amdahl, Gene M. 1967. "Validity of the Single Processor Approach to Achieving Large Scale Computing Capabilities." *Proceedings of the April 18-20, 1967, Spring Joint Computer Conference on - AFIPS '67 (Spring)*, AFIPS '67 (spring), 483–85. <https://doi.org/10.1145/1465482.1465560>.
- Amershi, Saleema, Andrew Begel, Christian Bird, Robert DeLine, Harald Gall, Ece Kamar, Nachiappan Nagappan, Besmira Nushi, and Thomas Zimmermann. 2019. "Software Engineering for Machine Learning: A Case Study." *2019 IEEE/ACM 41st International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP)*, 291–300. <https://doi.org/10.1109/icse-seip.2019.00042>.
- Amodei, Dario, and Danny Hernandez. 2018a. "AI and Compute." *OpenAI Blog* 2.
- Amodei, Dario, and Danny Hernandez. 2018b. "AI and Compute." *OpenAI Blog* 6.
- Angwin, Julia, Jeff Larson, Surya Mattu, and Lauren Kirchner. 2016. *Machine Bias*. ProPublica investigation.
- Ansel, Jason, Edward Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky, Bin Bao, et al. 2024. "PyTorch 2: Faster Machine Learning Through Dynamic Python Bytecode Transformation and Graph Compilation." *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2*, 929–47. <https://doi.org/10.1145/3620665.3640366>.
- Anthropic. 2023. *Introducing Claude 2.1*.

- Anthropic. 2024. *Introducing the Next Generation of Claude*.
- AnyLogic. 2024. *Synthetic Data for Artificial Intelligence*.
- Apache Software Foundation. 2024. *Apache Airflow*.
- Apple. 2024a. *Facial Matching Security*. Apple Platform Security.
- Apple. 2024b. *The Secure Enclave*. Apple Platform Security.
- Ardila, Rosana, Megan Branson, Kelly Davis, Michael Kohler, Josh Meyer, Michael Henretty, Reuben Morais, Lindsay Saunders, Francis Tyers, and Gregor Weber. 2020. "Common Voice: A Massively-Multilingual Speech Corpus." *Proceedings of the Twelfth Language Resources and Evaluation Conference*, 4218–22.
- Armbrust, Michael, Tathagata Das, Liwen Sun, Burak Yavuz, Shixiong Zhu, Mukul Murthy, Joseph Torres, et al. 2020. "Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores." *Proceedings of the VLDB Endowment* 13 (12): 3411–24. <https://doi.org/10.14778/3415478.3415560>.
- Authors, Kubeflow. 2024. *Kubeflow*.
- Ba, Jimmy Lei, Jamie Ryan Kiros, and Geoffrey E. Hinton. 2016. "Layer Normalization." *arXiv Preprint arXiv:1607.06450*.
- Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio. 2015. "Neural Machine Translation by Jointly Learning to Align and Translate." *International Conference on Learning Representations (ICLR)*.
- Bai, Zhaojun, James Demmel, Jack Dongarra, Julien Langou, and Jenny Wang. 2006. "LAPACK." In *Handbook of Linear Algebra*. Chapman; Hall/CRC. <https://doi.org/10.1201/9781420010572-75>.
- Baidu Research. 2016. *DeepBench: Benchmarking Deep Learning Operations on Different Hardware*.
- Bakheet, Samy, and Ayoub Al-Hamadi. 2021. "A Framework for Instantaneous Driver Drowsiness Detection Based on Improved HOG Features and Naïve Bayesian Classification." *Brain Sciences* 11 (2): 240. <https://doi.org/10.3390/brainsci11020240>.
- Banbury, Colby R., Vijay Janapa Reddi, Max Lam, William Fu, Amin Fazel, Jeremy Holleman, Xinyuan Huang, et al. 2020. "Benchmarking TinyML Systems: Challenges and Direction." *arXiv Preprint arXiv:2003.04821*.
- Banbury, Colby R, Chuteng Zhou, Igor Fedorov, Ramon Matas, Urmish Thakker, Dibakar Gope, Vijay Janapa Reddi, Matthew Mattina, and Paul N Whatmough. 2021. "MicroNets: Neural Network Architectures for Deploying TinyML Applications on Commodity Microcontrollers." *Proceedings of Machine Learning and Systems* 3: 517–32.
- Banbury, Colby, Vijay Janapa Reddi, Peter Torelli, Jeremy Holleman, Nat Jeffries, Csaba Kiraly, Pietro Montino, et al. 2021. "MLPerf Tiny Benchmark." *arXiv Preprint*.
- Barocas, Solon, and Andrew D. Selbst. 2016. "Big Data's Disparate Impact." *California Law Review* 104: 671–732. <https://doi.org/10.2139/ssrn.2477899>.
- Barroso, Luiz André, Urs Hölzle, and Parthasarathy Ranganathan. 2019. *The Datacenter as a Computer: Designing Warehouse-Scale Machines*. Synthesis Lectures on Computer Architecture. Springer International Publishing. <https://doi.org/10.1007/978-3-031-01761-2>.
- Barroso, Luiz, Mike Marty, David Patterson, and Parthasarathy Ranganathan. 2017. "Attack of the Killer Microseconds." *Communications of the ACM* 60 (4): 48–54. <https://doi.org/10.1145/3015146>.
- Baur, Walter, and Volker Strassen. 1983. "The Complexity of Partial Derivatives." *Theoretical Computer Science* 22 (3): 317–30. [https://doi.org/10.1016/0304-3975\(83\)90110-x](https://doi.org/10.1016/0304-3975(83)90110-x).
- Baydin, Atilim Gunes, Barak A. Pearlmutter, Alexey Andreyevich Radul, and Jeffrey Mark Siskind. 2018. "Automatic Differentiation in Machine Learning: A Survey." *Journal of Machine Learning Research* 18 (153): 1–43.
- Baylor, Denis, Eric Breck, Heng-Tze Cheng, Noah Fiedel, Chuan Yu Foo, Zakaria Haque, Salem Haykal, et al. 2017. "TFX: A TensorFlow-Based Production-Scale Machine Learning Platform." *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1387–95. <https://doi.org/10.1145/3097983.3098021>.
- Bedford Taylor, Michael. 2017. "The Evolution of Bitcoin Hardware." *Computer* 50 (9): 58–66. <https://doi.org/10.1109/mc.2017.3571056>.
- Beede, Emma, Elizabeth Baylor, Fred Hersch, Anna Iurchenko, Lauren Wilcox, Paisan Ruamviboonsuk, and Laura M. Vardoulakis. 2020. "A Human-Centered Evaluation of a Deep Learning System Deployed in Clinics for the Detection of Diabetic Retinopathy." *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI)*, 1–12. <https://doi.org/10.1145/3313831.3376718>.
- Belkin, Mikhail, Daniel Hsu, Siyuan Ma, and Soumik Mandal. 2019. "Reconciling Modern Machine-Learning Practice and the Classical Bias–Variance Trade-Off." *Proceedings of the National Academy of Sciences* 116 (32): 15849–54. <https://doi.org/10.1073/pnas.1903070116>.

- Bellamy, Rachel K. E., Kuntal Dey, Michael Hind, Seung Hoffman, Sylvain Houde, Karthikeyan Kannan, Pradeep Lohia, et al. 2019. "AI Fairness 360: An Extensible Toolkit for Detecting and Mitigating Algorithmic Bias." *IBM Journal of Research and Development* 63 (4/5): 4:1–15. <https://doi.org/10.1147/jrd.2019.2942287>.
- Bender, Emily M., and Batya Friedman. 2018. "Data Statements for Natural Language Processing: Toward Mitigating System Bias and Enabling Better Science." *Transactions of the Association for Computational Linguistics* 6: 587–604. https://doi.org/10.1162/tacl_a_00041.
- Bengio, Emmanuel, Pierre-Luc Bacon, Joelle Pineau, and Doina Precup. 2015. "Conditional Computation in Neural Networks for Faster Models." *arXiv Preprint arXiv:1511.06297*.
- Bengio, Yoshua, Aaron Courville, and Pascal Vincent. 2013. "Representation Learning: A Review and New Perspectives." *IEEE Transactions on Pattern Analysis and Machine Intelligence* 35 (8): 1798–828. <https://doi.org/10.1109/tpami.2013.50>.
- Bengio, Yoshua, Nicholas Léonard, and Aaron Courville. 2013. "Estimating or Propagating Gradients Through Stochastic Neurons for Conditional Computation." *arXiv Preprint arXiv:1308.3432*.
- Bengio, Yoshua, Jérôme Louradour, Ronan Collobert, and Jason Weston. 2009. "Curriculum Learning." *Proceedings of the 26th Annual International Conference on Machine Learning*, 41–48. <https://doi.org/10.1145/1553374.1553380>.
- Berger, Vance W., and YanYan Zhou. 2014. "Kolmogorov–Smirnov Test: Overview." In *Wiley StatsRef: Statistics Reference Online*. Wiley. <https://doi.org/10.1002/9781118445112.stat06558>.
- Bergstra, James, Olivier Breuleux, Frédéric Bastien, Pascal Lamblin, Razvan Pascanu, Guillaume Desjardins, Joseph Turian, David Warde-Farley, and Yoshua Bengio. 2010. "Theano: A CPU and GPU Math Compiler in Python." *Proceedings of the Python in Science Conference* 4: 18–24. <https://doi.org/10.25080/majora-92bf1922-003>.
- Beyer, Betsy, Chris Jones, Jennifer Peto, and Niall Richard Murphy. 2016. *Site Reliability Engineering: How Google Runs Production Systems*. O'Reilly Media.
- Beyer, Lucas, Olivier J. Hénaff, Alexander Kolesnikov, Xiaohua Zhai, and Aäron van den Oord. 2020. "Are We Done with ImageNet?" *arXiv Preprint arXiv:2006.07159*.
- Bird, Sarah, Miro Dudík, Richard Edgar, Brandon Horn, Roman Lutz, Vanessa Milan, Mehrnoosh Sameki, Hanna Wallach, and Kathleen Walker. 2020. "Fairlearn: A Toolkit for Assessing and Improving Fairness in AI." *Microsoft Technical Report MSR-TR-2020-32*.
- Bishop, Christopher M. 2006. *Pattern Recognition and Machine Learning*. Springer.
- Blalock, Davis, Jose Javier Gonzalez Ortiz, Jonathan Frankle, and John Gutttag. 2020. "What Is the State of Neural Network Pruning?" *Proceedings of Machine Learning and Systems (MLSys)* 2: 129–46.
- Bobrow, Daniel G. 1964. "Natural Language Input for a Computer Problem Solving System." PhD thesis, Massachusetts Institute of Technology.
- Boehm, Barry W. 1981. *Software Engineering Economics*. Prentice-Hall.
- Bolukbasi, Tolga, Kai-Wei Chang, James Y. Zou, Venkatesh Saligrama, and Adam Tauman Kalai. 2016. "Man Is to Computer Programmer as Woman Is to Homemaker? Debiasing Word Embeddings." *Advances in Neural Information Processing Systems (NeurIPS)*, 4349–57.
- Bommasani, Rishi, Drew A. Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S. Bernstein, et al. 2021. "On the Opportunities and Risks of Foundation Models." *arXiv Preprint arXiv:2108.07258*.
- Boroumand, Amirali, Saugata Ghose, Youngsok Kim, Rachata Ausavarungnirun, Eric Shiu, Rahul Thakur, Daehyun Kim, et al. 2018. "Google Workloads for Consumer Devices: Mitigating Data Movement Bottlenecks." *Proceedings of the Twenty-Third International Conference on Architectural Support for Programming Languages and Operating Systems, ASPLOS '18*, 316–31. <https://doi.org/10.1145/3173162.3173177>.
- Bourtole, Lucas, Varun Chandrasekaran, Christopher A. Choquette-Choo, Hengrui Jia, Adelin Travers, Baiwu Zhang, David Lie, and Nicolas Papernot. 2021. "Machine Unlearning." *2021 IEEE Symposium on Security and Privacy (SP)*, 141–59. <https://doi.org/10.1109/sp40001.2021.00019>.
- Bradbury, James, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, George Neca, et al. 2018. *JAX: Composable Transformations of Python+NumPy Programs*.
- Breck, Eric, Shanying Cai, Eric Nielsen, Michael Salib, and D. Sculley. 2017. "The ML Test Score: A Rubric for ML Production Readiness and Technical Debt Reduction." *Proceedings of the 2017 IEEE International Conference on Big Data (Big Data)*, 1123–32. <https://doi.org/10.1109/bigdata.2017.8258038>.
- Breck, Eric, Neoklis Polyzotis, Sudip Roy, Steven Euijong Whang, and Martin Zinkevich. 2019. "Data Validation for Machine Learning." *Proceedings of the 2nd SysML Conference* 1: 334–47.
- Brewer, Eric A. 2000. "Towards Robust Distributed Systems (Abstract)." *Proceedings of the Nineteenth Annual ACM Symposium on Principles of Distributed Computing*, 7. <https://doi.org/10.1145/343477.343502>.

- Broder, A. Z. 1997. "On the Resemblance and Containment of Documents." *Proceedings. Compression and Complexity of SEQUENCES 1997 (Cat. No.97TB100171)*, 21–29. <https://doi.org/10.1109/sequen.1997.666900>.
- Brown, Tom B., Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. 2020. "Language Models Are Few-Shot Learners." *Advances in Neural Information Processing Systems* 33: 1877–901. <https://doi.org/10.48550/arxiv.2005.14165>.
- Brutlag, Jake. 2009. *Speed Matters for Google Web Search*. Google Research Blog.
- Buolamwini, Joy, and Timnit Gebru. 2018. "Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification." *Conference on Fairness, Accountability and Transparency*, 77–91.
- Burns, Brendan, Brian Grant, David Oppenheimer, Eric Brewer, and John Wilkes. 2016. "Borg, Omega, and Kubernetes." *Communications of the ACM* 59 (5): 50–57. <https://doi.org/10.1145/2890784>.
- Cai, Han, Chuang Gan, Ligeng Zhu, and Song Han. 2020. "TinyTL: Reduce Activations, Not Trainable Parameters for Efficient on-Device Learning." *Advances in Neural Information Processing Systems* 33: 11285–97.
- Campbell, Murray, Jr. Hoane A. Joseph, and Feng-hsiung Hsu. 2002. "Deep Blue." *Artificial Intelligence* 134 (1-2): 57–83. [https://doi.org/10.1016/S0004-3702\(01\)00129-1](https://doi.org/10.1016/S0004-3702(01)00129-1).
- Cao, Yinzhi, and Junfeng Yang. 2015. "Towards Making Systems Forget with Machine Unlearning." *2015 IEEE Symposium on Security and Privacy*, 463–80. <https://doi.org/10.1109/sp.2015.35>.
- Chapman, Pete, Julian Clinton, Randy Kerber, Thomas Khabaza, Thomas Reinartz, Colin Shearer, and Rudiger Wirth. 2000. "CRISP-DM 1.0: Step-by-Step Data Mining Guide." *SPSS Inc* 1: 78.
- Chen, Andrew, Andy Chow, Aaron Davidson, Arjun DCunha, Ali Ghodsi, Sue Ann Hong, Andy Konwinski, et al. 2020. "Developments in MLflow: A System to Accelerate the Machine Learning Lifecycle." *Proceedings of the Fourth International Workshop on Data Management for End-to-End Machine Learning*, 1–4. <https://doi.org/10.1145/3399579.3399867>.
- Chen, Emma, Shvetank Prakash, Vijay Janapa Reddi, David Kim, and Pranav Rajpurkar. 2023. "A Framework for Integrating Artificial Intelligence for Clinical Care with Continuous Therapeutic Monitoring." *Nature Biomedical Engineering* 9 (4): 445–54. <https://doi.org/10.1038/s41551-023-01115-0>.
- Chen, Mark, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, et al. 2021. "Evaluating Large Language Models Trained on Code." *arXiv Preprint arXiv:2107.03374*.
- Chen, Tianqi, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Q. Yan, Haichen Shen, Meghan Cowan, et al. 2018. "TVM: An Automated End-to-End Optimizing Compiler for Deep Learning." *Proceedings of the 13th USENIX Symposium on Operating Systems Design and Implementation (OSDI '18)*, 578–94.
- Chen, Tianqi, Bing Xu, Chiyuan Zhang, and Carlos Guestrin. 2016. "Training Deep Nets with Sublinear Memory Cost." *arXiv Preprint arXiv:1604.06174*.
- Chen, Tianshi, Zidong Du, Ninghui Sun, Jia Wang, Chengyong Wu, Yunji Chen, and Olivier Temam. 2014. "Dian-Nao: A Small-Footprint High-Throughput Accelerator for Ubiquitous Machine-Learning." *Proceedings of the 19th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS '14)*, 269–84. <https://doi.org/10.1145/2541940.2541967>.
- Chen, Ting, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. 2020. "A Simple Framework for Contrastive Learning of Visual Representations." *Proceedings of the 37th International Conference on Machine Learning, ICML'20*, vol. 119: 1597–607.
- Chen, Yanxi, Xuchen Pan, Yaliang Li, Bolin Ding, and Jingren Zhou. 2024. "EE-LLM: Large-Scale Training and Inference of Early-Exit Large Language Models with 3D Parallelism." *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- Chen, Yu-Hsin, Joel S. Emer, and Vivienne Sze. 2016. "Eyeriss: A Spatial Architecture for Energy-Efficient Dataflow for Convolutional Neural Networks." *2016 ACM/IEEE 43rd Annual International Symposium on Computer Architecture (ISCA)*, 367–79. <https://doi.org/10.1109/ISCA.2016.40>.
- Chen, Yu-Hsin, Joel Emer, and Vivienne Sze. 2017. "Using Dataflow to Optimize Energy Efficiency of Deep Neural Network Accelerators." *IEEE Micro* 37 (3): 12–21. <https://doi.org/10.1109/mm.2017.54>.
- Chetlur, Sharan, Cliff Woolley, Philippe Vandermersch, Jonathan Cohen, John Tran, Bryan Catanzaro, and Evan Shelhamer. 2014. "cuDNN: Efficient Primitives for Deep Learning." *arXiv Preprint arXiv:1410.0759*.
- Cho, Aeree, Grace C. Kim, Alexander Karpekov, Alec Helbling, Zijie J. Wang, Seongmin Lee, Benjamin Hoover, and Duen Horng (Polo) Chau. 2025. "TRANSFORMER EXPLAINER: Interactive Learning of Text-Generative Models." *Proceedings of the AAAI Conference on Artificial Intelligence* 39 (28): 29625–27. <https://doi.org/10.1609/aaai.v39i28.35347>.
- Cho, Kyunghyun, Bart van Merriënboer, Dzmitry Bahdanau, and Yoshua Bengio. 2014. "On the Properties of Neural Machine Translation: Encoder-Decoder Approaches." *Proceedings of SSST-8, Eighth Workshop on Syntax, Semantics and Structure in Statistical Translation*, 103–11. <https://doi.org/10.3115/v1/w14-4012>.

- Choi, Dami, Alexandre Passos, Christopher J. Shallue, and George E. Dahl. 2019. "Faster Neural Network Training with Data Echoing." *arXiv Preprint*, ahead of print. <https://doi.org/10.48550/arXiv.1907.05550>.
- Choi, Jungwook, Zhuo Wang, Swagath Venkataramani, Pierce I-Jen Chuang, Vijayalakshmi Srinivasan, and Kailash Gopalakrishnan. 2018. "PACT: Parameterized Clipping Activation for Quantized Neural Networks." *arXiv Preprint*.
- Chollet, François. 2018. *Deep Learning with Python*. Manning Publications.
- Chollet, François et al. 2024. *Keras*.
- Choquette, Jack. 2023. "NVIDIA Hopper H100 GPU: Scaling Performance." *IEEE Micro* 43 (3): 9–17. <https://doi.org/10.1109/mm.2023.3256796>.
- Choquette, Jack, Wishwesh Gandhi, Olivier Giroux, Nick Stam, and Ronny Krashinsky. 2021. "NVIDIA A100 Tensor Core GPU: Performance and Innovation." *IEEE Micro* 41 (2): 29–35. <https://doi.org/10.1109/mm.2021.3061394>.
- Choudhary, Tejalal, Vipul Mishra, Anurag Goswami, and Jagannathan Sarangapani. 2020. "A Comprehensive Survey on Model Compression and Acceleration." *Artificial Intelligence Review* 53 (7): 5113–55. <https://doi.org/10.1007/s10462-020-09816-7>.
- Choukroun, Yoni, Eli Kravchik, Fan Yang, and Pavel Kisilev. 2019. "Low-Bit Quantization of Neural Networks for Efficient Inference." *2019 IEEE/CVF International Conference on Computer Vision Workshop (ICCVW)*, 3009–18. <https://doi.org/10.1109/iccvw.2019.00363>.
- Chouldechova, Alexandra. 2017. "Fair Prediction with Disparate Impact: A Study of Bias in Recidivism Prediction Instruments." *Big Data* 5 (2): 153–63. <https://doi.org/10.1089/big.2016.0047>.
- Chowdhery, Aakanksha, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, et al. 2022. "PaLM: Scaling Language Modeling with Pathways." *arXiv Preprint arXiv:2204.02311*.
- Chu, Grace, Okan Arikan, Gabriel Bender, Weijun Wang, Achille Brighton, Pieter-Jan Kindermans, Hanxiao Liu, Berkin Akin, Suyog Gupta, and Andrew Howard. 2021. "Discovering Multi-Hardware Mobile Models via Architecture Search." *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 3016–25. <https://doi.org/10.1109/cvprw53098.2021.00337>.
- CircleCI. 2024. *CircleCI: Continuous Integration and Delivery Platform*.
- Clark, Kevin, Urvashi Khandelwal, Omer Levy, and Christopher D. Manning. 2019. "What Does BERT Look at? An Analysis of BERT's Attention." *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, 276–86. <https://doi.org/10.18653/v1/w19-4828>.
- Cloud Native Computing Foundation. 2024a. *Kubernetes: Production-Grade Container Orchestration*.
- Cloud Native Computing Foundation. 2024b. *Prometheus: Monitoring System and Time Series Database*.
- Cohen, Jacob. 1960. "A Coefficient of Agreement for Nominal Scales." *Educational and Psychological Measurement* 20 (1): 37–46. <https://doi.org/10.1177/001316446002000104>.
- Cohen, Taco, and Max Welling. 2016. "Group Equivariant Convolutional Networks." *Proceedings of the 33rd International Conference on Machine Learning (ICML)* 48: 2990–99.
- Coleman, Cody, Edward Chou, Julian Katz-Samuels, Sean Culatana, Peter Bailis, Alexander C. Berg, Robert Nowak, Roshan Sumbaly, Matei Zaharia, and I. Zeki Yalniz. 2022. "Similarity Search for Efficient Active Learning and Search of Rare Concepts." *Proceedings of the AAAI Conference on Artificial Intelligence* 36 (6): 6402–10. <https://doi.org/10.1609/aaai.v36i6.20591>.
- Coleman, Cody, Deepak Narayanan, Daniel Kang, Tian Zhao, Jian Zhang, Luigi Nardi, Peter Bailis, Kunle Olukotun, Christopher Ré, and Matei Zaharia. 2019. "Analysis of DAWNbench, a Time-to-Accuracy Machine Learning Performance Benchmark." *ACM SIGOPS Operating Systems Review* 53 (1): 14–25. <https://doi.org/10.1145/3352020.3352024>.
- Courbariaux, Matthieu, Yoshua Bengio, and Jean-Pierre David. 2015. "BinaryConnect: Training Deep Neural Networks with Binary Weights During Propagations." *Advances in Neural Information Processing Systems (NeurIPS)* 28: 3123–31.
- Cover, Thomas M., and Joy A. Thomas. 2006. *Elements of Information Theory*. 2nd ed. Wiley. <https://doi.org/10.1002/047174882X>.
- Covington, Paul, Jay Adams, and Emre Sargin. 2016. "Deep Neural Networks for YouTube Recommendations." *Proceedings of the 10th ACM Conference on Recommender Systems*, 191–98. <https://doi.org/10.1145/2959100.2959190>.
- Crankshaw, Daniel, Xin Wang, Guilio Zhou, Michael J. Franklin, Joseph E. Gonzalez, and Ion Stoica. 2017. "Clipper: A Low-Latency Online Prediction Serving System." *14th USENIX Symposium on Networked Systems Design and Implementation (NSDI 17)*, 613–27.
- CrowdFlower. 2016. *Data Science Report 2016*.
- Cubuk, Ekin D., Barret Zoph, Dandelion Mané, Vijay Vasudevan, and Quoc V. Le. 2019. "AutoAugment: Learning Augmentation Strategies from Data." *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 113–23. <https://doi.org/10.1109/cvpr.2019.00020>.

- Cubuk, Ekin D., Barret Zoph, Jonathon Shlens, and Quoc V. Le. 2020. "Randaugment: Practical Automated Data Augmentation with a Reduced Search Space." *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 3008–17. <https://doi.org/10.1109/cvprw50498.2020.00359>.
- Cunningham, Ward. 1992. "The WyCash Portfolio Management System." *ACM SIGPLAN OOPS Messenger* 4 (2): 29–30. <https://doi.org/10.1145/157710.157715>.
- Curnow, H. J., and B. A. Wichmann. 1976. "A Synthetic Benchmark." *The Computer Journal* 19 (1): 43–49. <https://doi.org/10.1093/comjnl/19.1.43>.
- Cybenko, George. 1989. "Approximation by Superpositions of a Sigmoidal Function." *Mathematics of Control, Signals, and Systems* 2 (4): 303–14. <https://doi.org/10.1007/bf02551274>.
- Dalal, Navneet, and Bill Triggs. 2005. "Histograms of Oriented Gradients for Human Detection." *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)* 1: 886–93. <https://doi.org/10.1109/cvpr.2005.177>.
- Dally, Bill. 2023. "Hardware for Deep Learning." *2023 IEEE Hot Chips 35 Symposium (HCS)*, 1–58. <https://doi.org/10.1109/hcs59251.2023.10254716>.
- Dally, William J., Stephen W. Keckler, and David B. Kirk. 2021. "Evolution of the Graphics Processing Unit (GPU)." *IEEE Micro* 41 (6): 42–51. <https://doi.org/10.1109/mm.2021.3113475>.
- Dao, T., D. Y. Fu, S. Ermon, A. Rudra, and C. Ré. 2022. "FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness." *Advances in Neural Information Processing Systems (NeurIPS)* 35: 16344–59. <https://doi.org/10.52202/068431-1189>.
- Dao, Tri. 2023. "FlashAttention-2: Faster Attention with Better Parallelism and Work Partitioning." *arXiv Preprint arXiv:2307.08691*.
- Dao, Tri, Beidi Chen, Nimit Sohoni, Arjun Desai, Michael Poli, Jessica Grogan, Alexander Liu, Aniruddh Rao, Atri Rudra, and Christopher Ré. 2022. "Monarch: Expressive Structured Matrices for Efficient and Accurate Training." *arXiv Preprint*.
- Dastin, Jeffrey. 2018. "Amazon Scraps Secret AI Recruiting Tool That Showed Bias Against Women." *Reuters*.
- Data Protection Commission. 2023. *Data Protection Commission Announces Conclusion of Two Inquiries into Meta Ireland*. Regulatory press release.
- Databricks. 2024. *MLflow: An Open Source Platform for the Machine Learning Lifecycle*.
- David, Robert, Jared Duke, Advait Jain, Vijay Janapa Reddi, Nat Jeffries, Jian Li, Nick Kreeger, et al. 2021. "TensorFlow Lite Micro: Embedded Machine Learning for TinyML Systems." *Proceedings of Machine Learning and Systems* 3: 800–811.
- dbt Labs. 2024. *Dbt (Data Build Tool)*.
- Dean, Jeff, David Patterson, and Cliff Young. 2018. "A New Golden Age in Computer Architecture: Empowering the Machine-Learning Revolution." *IEEE Micro* 38 (2): 21–29. <https://doi.org/10.1109/mm.2018.112130030>.
- Dean, Jeffrey. 2012. *Achieving Rapid Response Times in Large Online Services*. Berkeley AMPLab Cloud Seminar.
- Dean, Jeffrey, and Luiz André Barroso. 2013. "The Tail at Scale." *Communications of the ACM* 56 (2): 74–80. <https://doi.org/10.1145/2408776.2408794>.
- Dean, Jeffrey, Greg Corrado, Rajat Monga, Kai Chen, Matthieu Devin, Quoc V. Le, Mark Z. Mao, et al. 2012. "Large Scale Distributed Deep Networks." In *Advances in Neural Information Processing Systems (NeurIPS)*, edited by Peter L. Bartlett, Fernando C. N. Pereira, Christopher J. C. Burges, Léon Bottou, and Kilian Q. Weinberger, vol. 25. Curran Associates.
- Dean, Jeffrey, and Sanjay Ghemawat. 2004. "MapReduce: Simplified Data Processing on Large Clusters." *Proceedings of the 6th Symposium on Operating Systems Design and Implementation (OSDI)*, 137–50.
- DeepSeek. 2024. *Introducing DeepSeek-V3*.
- Dehghani, Zhamak. 2022. *Data Mesh: Delivering Data-Driven Value at Scale*. O'Reilly Media.
- Deng, Jia, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. "ImageNet: A Large-Scale Hierarchical Image Database." *2009 IEEE Conference on Computer Vision and Pattern Recognition*, 248–58. <https://doi.org/10.1109/cvpr.2009.5206848>.
- Deng, Jia, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2024. *ImageNet*.
- Dennard, Robert H., Frank H. Gaensslen, Hwa-Nien Yu, Victor L. Rideout, Elias Bassous, and Antoine R. LeBlanc. 1974. "Design of Ion-Implanted MOSFET's with Very Small Physical Dimensions." *IEEE J. Solid-State Circuits* 9 (5): 256–68. <https://doi.org/10.1109/jssc.1974.1050511>.
- Denning, Peter J. 1968. "The Working Set Model for Program Behavior." *Communications of the ACM* 11 (5): 323–33. <https://doi.org/10.1145/363095.363141>.

- Dettmers, Tim, Mike Lewis, Sam Shleifer, and Luke Zettlemoyer. 2022. "8-Bit Optimizers via Block-Wise Quantization." *International Conference on Learning Representations (ICLR)*.
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. "BERT: Pre-Training of Deep Bidirectional Transformers for Language Understanding." *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 4171–86. <https://doi.org/10.18653/v1/n19-1423>.
- DeVries, Terrance, and Graham W. Taylor. 2017. "Improved Regularization of Convolutional Neural Networks with Cutout." *arXiv Preprint arXiv:1708.04552*.
- Dixit, Kaivalya M. 1993. "Overview of the SPEC Benchmarks." In *The Benchmark Handbook for Database and Transaction Systems*, 2nd ed., edited by Jim Gray. Morgan Kaufmann.
- Dongarra, J. J., J. R. Bunch, C. B. Moler, and G. W. Stewart. 1979. *LINPACK Users' Guide*. SIAM. <https://doi.org/10.1137/1.9781611971811>.
- Dongarra, Jack J., Jeremy Du Croz, Sven Hammarling, and Richard J. Hanson. 1988. "An Extended Set of FORTRAN Basic Linear Algebra Subprograms." *ACM Transactions on Mathematical Software* 14 (1): 1–17. <https://doi.org/10.1145/42288.42291>.
- Dosovitskiy, Alexey, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, et al. 2021. "An Image Is Worth 16x16 Words: Transformers for Image Recognition at Scale." *International Conference on Learning Representations (ICLR)*.
- Dua, Dheeru, and Casey Graff. 2024. *UCI Machine Learning Repository*.
- Duarte, Javier, Nhan Tran, Ben Hawks, Christian Herwig, Jules Muhizi, Shvetank Prakash, and Vijay Janapa Reddi. 2022. "FastML Science Benchmarks: Accelerating Real-Time Scientific Edge Machine Learning." *arXiv Preprint arXiv:2207.07958*.
- Dubey, Abhimanyu, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, et al. 2024. *The Llama 3 Herd of Models*. arXiv preprint arXiv:2407.21783.
- Dwork, Cynthia. 2008. "Differential Privacy: A Survey of Results." In *Theory and Applications of Models of Computation*. Springer Berlin Heidelberg. https://doi.org/10.1007/978-3-540-79228-4_1.
- Dwork, Cynthia, Frank McSherry, Kobbi Nissim, and Adam Smith. 2006. "Calibrating Noise to Sensitivity in Private Data Analysis." In *Theory of Cryptography Conference (TCC)*, edited by Shai Halevi and Tal Rabin, vol. 3876. Lecture Notes in Computer Science. Springer Berlin Heidelberg. https://doi.org/10.1007/11681878_14.
- Elastic NV. 2024. *Elasticsearch: Distributed Search and Analytics Engine*.
- Elman, Jeffrey L. 1990. "Finding Structure in Time." *Cognitive Science* 14 (2): 179–211. https://doi.org/10.1207/s15516709cog1402_1.
- Elsen, Erich, Marat Dukhan, Trevor Gale, and Karen Simonyan. 2020. "Fast Sparse ConvNets." *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 14617–26. <https://doi.org/10.1109/cvpr42600.2020.01464>.
- Elsken, Thomas, Jan Hendrik Metzen, and Frank Hutter. 2019. "Neural Architecture Search." In *The Springer Series on Challenges in Machine Learning*, vol. 20. Springer International Publishing. https://doi.org/10.1007/978-3-030-05318-5_3.
- Engineering, M. 2016. *Introducing FBLearner Flow: Facebook's AI Backbone*. Engineering at Meta Blog.
- Epoch AI. 2024. *Machine Learning Trends*. Epoch AI Research Database.
- Equal Employment Opportunity Commission, Civil Service Commission, Department of Labor, and Department of Justice. 1978. *Uniform Guidelines on Employee Selection Procedures*. Code of Federal Regulations.
- Esmailzadeh, Hadi, Emily Blem, Renee St. Amant, Karthikeyan Sankaralingam, and Doug Burger. 2011. "Dark Silicon and the End of Multicore Scaling." *Proceedings of the 38th Annual International Symposium on Computer Architecture*, 365–76. <https://doi.org/10.1145/2000064.2000108>.
- Ethernet Alliance. 2025. *2025 Ethernet Roadmap*. Ethernet Alliance roadmap.
- European Data Protection Board. 2018. *Guidelines on Automated Individual Decision-Making and Profiling for the Purposes of Regulation 2016/679*.
- European Parliament and Council of the European Union. 2024. *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying down Harmonised Rules on Artificial Intelligence (AI Act)*. Official Journal of the European Union, L 2024/1689.
- European Parliament, and Council of the European Union. 2016. *Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016*. Official Journal of the European Union.
- European Space Agency. 1996. *Ariane 501: Presentation of Inquiry Board Report*. European Space Agency press release.
- Everingham, Mark, Luc Van Gool, Christopher K. I. Williams, John Winn, and Andrew Zisserman. 2009. "The Pascal Visual Object Classes (VOC) Challenge." *International Journal of Computer Vision* 88 (2): 303–38. <https://doi.org/10.1007/s11263-009-0275-4>.

- Fedus, William, Barret Zoph, and Noam Shazeer. 2022. "Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity." *Journal of Machine Learning Research* 23 (120): 1–39.
- Feigenbaum, Edward A. 1984. "Knowledge Engineering: The Applied Side of Artificial Intelligence." *Annals of the New York Academy of Sciences* 426: 91–107. <https://doi.org/10.1111/j.1749-6632.1984.tb16513.x>.
- Feng, Wu-chun, and Kirk W. Cameron. 2007. "The Green500 List: Encouraging Sustainable Supercomputing." *Computer* 40 (12): 50–55. <https://doi.org/10.1109/MC.2007.445>.
- Fitzpatrick, Thomas B. 1988. "The Validity and Practicality of Sun-Reactive Skin Types i Through VI." *Archives of Dermatology* 124 (6): 869. <https://doi.org/10.1001/archderm.1988.01670060015008>.
- FlatBuffers Authors. 2026. *FlatBuffers Documentation*.
- Fleiss, Joseph L. 1971. "Measuring Nominal Scale Agreement Among Many Raters." *Psychological Bulletin* 76 (5): 378–82. <https://doi.org/10.1037/h0031619>.
- Flynn, M. J. 1966. "Very High-Speed Computing Systems." *Proceedings of the IEEE* 54 (12): 1901–9. <https://doi.org/10.1109/proc.1966.5273>.
- Fowler, Martin. 2014. *Canary Release*. Martin Fowler's Blog.
- Frankle, Jonathan, and Michael Carbin. 2019. "The Lottery Ticket Hypothesis: Finding Sparse, Trainable Neural Networks." *International Conference on Learning Representations (ICLR)*.
- Frostig, Roy, Matthew James Johnson, and Chris Leary. 2018. "Compiling Machine Learning Programs via High-Level Tracing." *Systems for Machine Learning*.
- Gabor, D. 1946. "Theory of Communication. Part 1: The Analysis of Information." *Journal of the Institution of Electrical Engineers - Part III: Radio and Communication Engineering* 93 (26): 429–41. <https://doi.org/10.1049/ji-3-2.1946.0074>.
- Gadre, Samir Yitzhak, Gabriel Ilharco, Alex Fang, Jonathan Hayase, Georgios Smyrnis, Thao Nguyen, Ryan Marten, et al. 2023. "DataComp: In Search of the Next Generation of Multimodal Datasets." *Advances in Neural Information Processing Systems (NeurIPS)* 36: 27092–112. <https://doi.org/10.52202/075280-1179>.
- Gale, Trevor, Erich Elsen, and Sara Hooker. 2019. "The State of Sparsity in Deep Neural Networks." *arXiv Preprint arXiv:1902.09574*.
- Gale, Trevor, Deepak Narayanan, Cliff Young, and Matei Zaharia. 2022. "MegaBlocks: Efficient Sparse Training with Mixture-of-Experts." *arXiv Preprint*.
- Gale, Trevor, Matei Zaharia, Cliff Young, and Erich Elsen. 2020. "Sparse GPU Kernels for Deep Learning." *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis* 2020: 8955–67. <https://doi.org/10.1109/sc41405.2020.00021>.
- Gama, João, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. 2014. "A Survey on Concept Drift Adaptation." *ACM Computing Surveys* 46 (4): 1–37. <https://doi.org/10.1145/2523813>.
- Gao, Leo, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, et al. 2020. "The Pile: An 800GB Dataset of Diverse Text for Language Modeling." *arXiv Preprint arXiv:2101.00027*.
- Gartner. 2024. *Gartner Forecasts Worldwide Public Cloud End-User Spending to Total \$679 Billion in 2024*. Gartner, Inc.
- Gebru, Timnit, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. 2021. "Datasheets for Datasets." *Communications of the ACM* 64 (12): 86–92. <https://doi.org/10.1145/3458723>.
- Gholami, Amir, Sehoon Kim, Zhen Dong, Zhewei Yao, Michael W. Mahoney, and Kurt Keutzer. 2021. "A Survey of Quantization Methods for Efficient Neural Network Inference." *arXiv Preprint arXiv:2103.13630* abs/2103.13630.
- Gholami, Amir, Sehoon Kim, Zhen Dong, Zhewei Yao, Michael W. Mahoney, and Kurt Keutzer. 2022. "A Survey of Quantization Methods for Efficient Neural Network Inference." In *Low-Power Computer Vision: Improve the Efficiency of Artificial Intelligence*. Chapman; Hall/CRC. <https://doi.org/10.1201/9781003162810-13>.
- Gholami, Amir, Zhewei Yao, Sehoon Kim, Coleman Hooper, Michael W. Mahoney, and Kurt Keutzer. 2024. "AI and Memory Wall." *IEEE Micro* 44 (3): 33–39. <https://doi.org/10.1109/mm.2024.3373763>.
- Gilbert, Seth, and Nancy Lynch. 2002. "Brewer's Conjecture and the Feasibility of Consistent, Available, Partition-Tolerant Web Services." *ACM SIGACT News* 33 (2): 51–59. <https://doi.org/10.1145/564585.564601>.
- Ginsberg, Jeremy, Matthew H. Mohebbi, Rajan S. Patel, Lynnette Brammer, Mark S. Smolinski, and Larry Brilliant. 2009. "Detecting Influenza Epidemics Using Search Engine Query Data." *Nature* 457 (7232): 1012–14. <https://doi.org/10.1038/nature07634>.
- GitHub, Inc. 2024a. *GitHub*.
- GitHub, Inc. 2024b. *GitHub Actions*.

- Glorot, Xavier, and Yoshua Bengio. 2010. "Understanding the Difficulty of Training Deep Feedforward Neural Networks." *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics (AISTATS)* 9: 249–56.
- Gojek, and Google. 2019. *Feast: An Open Source Feature Store for Machine Learning*. Google Cloud Blog.
- Goodfellow, I. J., J. Shlens, and C. Szegedy. 2015. "Explaining and Harnessing Adversarial Examples." *International Conference on Learning Representations (ICLR)*.
- Goodfellow, Ian, Yoshua Bengio, and Aaron Courville. 2016. *Deep Learning*. MIT Press.
- Goodhart, Charles A. E. 1984. "Problems of Monetary Management: The UK Experience." In *Monetary Theory and Practice*. Palgrave. https://doi.org/10.1007/978-1-349-17295-5_4.
- Google. 2024a. *Our Next-Generation Model: Gemini 1.5*.
- Google. 2024b. *Static Vs. Dynamic Inference*. Google Machine Learning Crash Course.
- Google. 2025. *XLA: Optimizing Compiler for Machine Learning*.
- Google Cloud. 2024a. *Google Cloud Platform Documentation*. <https://cloud.google.com/docs>.
- Google Cloud. 2024b. *Google Cloud Storage*.
- Google Cloud. 2024c. *Tune Models Overview*.
- Google Cloud. 2024d. *Vertex AI Model Registry*.
- Google Cloud. 2026a. *Explore Models in Model Garden*.
- Google Cloud. 2026b. *MLOps: Continuous Delivery and Automation Pipelines in Machine Learning*.
- Google Cloud. 2026c. *Spot VMs*.
- Google Cloud. 2026d. *TPU v6e*.
- Google Developers Blog. 2024. *Gemini 1.5 Flash Price Drop with Tuning Rollout Complete, and More*.
- Gordon, Mitchell, Kevin Duh, and Nicholas Andrews. 2020. "Compressing BERT: Studying the Effects of Weight Pruning on Transfer Learning." *Proceedings of the 5th Workshop on Representation Learning for NLP*, 143–55. <https://doi.org/10.18653/v1/2020.repl4nlp-1.18>.
- Goto, Kazushige, and Robert A. van de Geijn. 2008. "Anatomy of High-Performance Matrix Multiplication." *ACM Transactions on Mathematical Software* 34 (3): 1–25. <https://doi.org/10.1145/1356052.1356053>.
- Gou, Jianping, Baosheng Yu, Stephen J. Maybank, and Dacheng Tao. 2021. "Knowledge Distillation: A Survey." *International Journal of Computer Vision* 129 (6): 1789–819. <https://doi.org/10.1007/s11263-021-01453-z>.
- Goyal, Priya, Piotr Dollár, Ross Girshick, Pieter Noordhuis, Lukasz Wesolowski, Aapo Kyrola, Andrew Tulloch, Yangqing Jia, and Kaiming He. 2017. "Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour." *arXiv Preprint arXiv:1706.02677* abs/1706.02677.
- Graphcore. 2020. *The Colossus MK2 IPU Processor*.
- Gray, Allison. 2015. *NVIDIA and IBM Cloud Support ImageNet Large Scale Visual Recognition Challenge*. NVIDIA Technical Blog.
- Gretton, Arthur, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola. 2012. "A Kernel Two-Sample Test." *Journal of Machine Learning Research* 13 (25): 723–73.
- Groeneveld, Dirk, Iz Beltagy, Pete Walsh, Akshita Bhagia, Rodney Kinney, Oyvind Tafjord, Ananya Harsh Jha, et al. 2024. "OLMo: Accelerating the Science of Language Models." *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 15789–809. <https://doi.org/10.18653/v1/2024.acl-long.841>.
- gRPC Authors. 2026. *Introduction to gRPC*.
- Gudivada, Venkat N., Amy Apon, and Junhua Ding. 2017. "Data Quality Considerations for Big Data and Machine Learning: Going Beyond Data Cleaning and Transformations." *International Journal on Advances in Software* 10 (1–2): 1–20.
- Gujarati, Arpan, Reza Karimi, Safya Alzayat, Wei Hao, Antoine Kaufmann, Ymir Vigfusson, and Jonathan Mace. 2020. "Clockwork: Predictable and Scalable DNN Inference in the Cloud." *USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 443–62.
- Gulshan, Varun, Lily Peng, Marc Coram, Martin C. Stumpe, Derek Wu, Arunachalam Narayanaswamy, Subhashini Venugopalan, et al. 2016. "Development and Validation of a Deep Learning Algorithm for Detection of Diabetic Retinopathy in Retinal Fundus Photographs." *JAMA* 316 (22): 2402. <https://doi.org/10.1001/jama.2016.17216>.

- Guo, Chuan, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. "On Calibration of Modern Neural Networks." *International Conference on Machine Learning (ICML)* 70: 1321–30.
- Gupta, Suyog, Ankur Agrawal, Kailash Gopalakrishnan, and Pritish Narayanan. 2015. "Deep Learning with Limited Numerical Precision." *International Conference on Machine Learning (ICML)*, 1737–46.
- Gupta, Udit, Carole-Jean Wu, Xiaodong Wang, Maxim Naumov, Brandon Reagen, David Brooks, Bradford Cotel, et al. 2020. "The Architectural Implications of Facebook's DNN-Based Personalized Recommendation." *2020 IEEE International Symposium on High Performance Computer Architecture (HPCA)*, 488–501. <https://doi.org/10.1109/HPCA47549.2020.00047>.
- Gustafson, John L. 1988. "Reevaluating Amdahl's Law." *Communications of the ACM* 31 (5): 532–33. <https://doi.org/10.1145/42411.42415>.
- Halevy, Alon, Peter Norvig, and Fernando Pereira. 2009. "The Unreasonable Effectiveness of Data." *IEEE Intelligent Systems* 24 (2): 8–12. <https://doi.org/10.1109/mis.2009.36>.
- Han, Song, Xingyu Liu, Huizi Mao, Jing Pu, Ardavan Pedram, Mark A. Horowitz, and William J. Dally. 2016. "EIE: Efficient Inference Engine on Compressed Deep Neural Network." *2016 ACM/IEEE 43rd Annual International Symposium on Computer Architecture (ISCA)*, 243–54. <https://doi.org/10.1109/isca.2016.30>.
- Han, Song, Huizi Mao, and William J. Dally. 2016. "Deep Compression: Compressing Deep Neural Networks with Pruning, Trained Quantization and Huffman Coding." *arXiv Preprint*.
- Han, Song, Jeff Pool, John Tran, and William J. Dally. 2015. "Learning Both Weights and Connections for Efficient Neural Networks." *Advances in Neural Information Processing Systems* 28 (*NeurIPS 2015*), 1135–43.
- Harchol-Balter, Mor. 2013. *Performance Modeling and Design of Computer Systems: Queueing Theory in Action*. Cambridge University Press. <https://doi.org/10.1017/cbo9781139226424>.
- Hardt, Moritz, Eric Price, and Nathan Srebro. 2016. "Equality of Opportunity in Supervised Learning." *Advances in Neural Information Processing Systems (NeurIPS)* 29: 3315–23.
- Harris, Charles R., K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, et al. 2020. "Array Programming with NumPy." *Nature* 585 (7825): 357–62. <https://doi.org/10.1038/s41586-020-2649-2>.
- Harvard Law School. 2024. *Does ChatGPT Violate New York Times Copyrights?*
- HashiCorp. 2014. *Terraform: Infrastructure as Code*. Software available from <https://www.terraform.io/>.
- Hatcher, Blake Douglas. 2024. "Automating Server Deployments with Ansible: Utilizing Automation in DevOps." *Journal of Computing Sciences in Colleges* 40 (3): 42–43.
- Hazelwood, Kim, Sarah Bird, David Brooks, Soumith Chintala, Utku Diril, Dmytro Dzhulgakov, Mohamed Fawzy, et al. 2018. "Applied Machine Learning at Facebook: A Datacenter Infrastructure Perspective." *2018 IEEE International Symposium on High Performance Computer Architecture (HPCA)*, 620–29. <https://doi.org/10.1109/hpca.2018.00059>.
- He, Kaiming, Haoqi Fan, Yuxin Wu, Saining Xie, and Ross Girshick. 2020. "Momentum Contrast for Unsupervised Visual Representation Learning." *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 9726–35. <https://doi.org/10.1109/cvpr42600.2020.00975>.
- He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2015. "Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification." *2015 IEEE International Conference on Computer Vision (ICCV)*, 1026–34. <https://doi.org/10.1109/iccv.2015.123>.
- He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016a. "Deep Residual Learning for Image Recognition." *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770–78. <https://doi.org/10.1109/cvpr.2016.90>.
- He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016b. "Identity Mappings in Deep Residual Networks." *Computer Vision – ECCV 2016*, 630–45. https://doi.org/10.1007/978-3-319-46493-0_38.
- Henderson, Peter, Jieru Hu, Joshua Romoff, Emma Brunskill, Dan Jurafsky, and Joelle Pineau. 2020. "Towards the Systematic Reporting of the Energy and Carbon Footprints of Machine Learning." *CoRR abs/2002.05651* (248): 1–43. <https://doi.org/10.48550/arxiv.2002.05651>.
- Henderson, Peter, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. 2018. "Deep Reinforcement Learning That Matters." *Proceedings of the AAAI Conference on Artificial Intelligence* 32 (1): 3207–14. <https://doi.org/10.1609/aaai.v32i1.11694>.
- Hendler, James A. 2008. "Avoiding Another AI Winter." *IEEE Intelligent Systems* 23 (2): 2–4. <https://doi.org/10.1109/MIS.2008.20>.
- Hendrycks, Dan, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2020. "Measuring Massive Multitask Language Understanding." *arXiv Preprint*.
- Hendrycks, Dan, and Kevin Gimpel. 2016. "Gaussian Error Linear Units (GELUs)." *arXiv Preprint arXiv:1606.08415*.
- Hennessy, J. L., and D. A. Patterson. 2011. *Computer Architecture: A Quantitative Approach*. Morgan Kaufmann.

- Hennessy, John L., and David A. Patterson. 2017. *Computer Architecture: A Quantitative Approach*. 6th ed. Morgan Kaufmann.
- Hennessy, John L., and David A. Patterson. 2019. "A New Golden Age for Computer Architecture." *Communications of the ACM* 62 (2): 48–60. <https://doi.org/10.1145/3282307>.
- Hermann, Jeremy, and Mike Del Balso. 2017. *Meet Michelangelo: Uber's Machine Learning Platform*. Uber Engineering Blog.
- Hernandez, Danny, and Tom B. Brown. 2020. "Measuring the Algorithmic Efficiency of Neural Networks." *arXiv Preprint arXiv:2005.04305*, ahead of print. <https://doi.org/10.48550/arxiv.2005.04305>.
- Hestness, Joel, Sharan Narang, Newsha Ardalani, Gregory Diamos, Heewoo Jun, Hassan Kianinejad, Md. Mostofa Ali Patwary, Yang Yang, and Yanqi Zhou. 2017. "Deep Learning Scaling Is Predictable, Empirically." *arXiv Preprint arXiv:1712.00409*.
- Hinton, Geoffrey, Oriol Vinyals, and Jeff Dean. 2015. "Distilling the Knowledge in a Neural Network." *arXiv Preprint*.
- Hirschberg, Julia, and Christopher D. Manning. 2015. "Advances in Natural Language Processing." *Science* 349 (6245): 261–66. <https://doi.org/10.1126/science.aaa8685>.
- Hochreiter, Sepp, and Jürgen Schmidhuber. 1997. "Long Short-Term Memory." *Neural Computation* 9 (8): 1735–80. <https://doi.org/10.1162/neco.1997.9.8.1735>.
- Hoefler, Torsten, Dan Alistarh, Tal Ben-Nun, Nikoli Dryden, and Alexandra Peste. 2021. "Sparsity in Deep Learning: Pruning and Growth for Efficient Inference and Training in Neural Networks." *Journal of Machine Learning Research* 22 (241): 1–124.
- Hoffmann, Jordan, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, et al. 2022. "Training Compute-Optimal Large Language Models." *Advances in Neural Information Processing Systems (NeurIPS)* 35: 30016–30. <https://doi.org/10.52202/068431-2176>.
- Holtzman, Ari, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. 2020. "The Curious Case of Neural Text Degeneration." *Proceedings of the 8th International Conference on Learning Representations (ICLR 2020)*.
- Hooker, Sara. 2021. "The Hardware Lottery." *Communications of the ACM* 64 (12): 58–65. <https://doi.org/10.1145/3467017>.
- Hornik, Kurt, Maxwell Stinchcombe, and Halbert White. 1989. "Multilayer Feedforward Networks Are Universal Approximators." *Neural Networks* 2 (5): 359–66. [https://doi.org/10.1016/0893-6080\(89\)90020-8](https://doi.org/10.1016/0893-6080(89)90020-8).
- Horowitz, Mark. 2014. "1.1 Computing's Energy Problem (and What We Can Do about It)." *2014 IEEE International Solid-State Circuits Conference Digest of Technical Papers (ISSCC)*, 10–14. <https://doi.org/10.1109/isscc.2014.6757323>.
- Howard, A. G., M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam. 2017. "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications." *CoRR* abs/1704.04861.
- Howard, Andrew, Mark Sandler, Bo Chen, Weijun Wang, Liang-Chieh Chen, Mingxing Tan, Grace Chu, et al. 2019. "Searching for MobileNetV3." *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, 1314–24. <https://doi.org/10.1109/iccv.2019.00140>.
- Howarth, Jesse. 2020. "Why Discard Is Switching from Go to Rust."
- Hu, Jie, Peng Lin, Huajun Zhang, Zining Lan, Wenxin Chen, Kailiang Xie, Siyun Chen, Hao Wang, and Sheng Chang. 2023. "A Dynamic Pruning Method on Multiple Sparse Structures in Deep Neural Networks." *IEEE Access* 11: 38448–57. <https://doi.org/10.1109/access.2023.3267469>.
- Hu, Ting-Kuei, Tianlong Chen, Haotao Wang, and Zhangyang Wang. 2020. "Triple Wins: Boosting Accuracy, Robustness and Efficiency Together by Enabling Input-Adaptive Inference." *International Conference on Learning Representations (ICLR)*.
- Huang, Gao, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q. Weinberger. 2017. "Densely Connected Convolutional Networks." *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2261–69. <https://doi.org/10.1109/cvpr.2017.243>.
- Huang, Y., Y. Cheng, A. Bapna, O. Firat, D. Chen, M. X. Chen, H. Lee, et al. 2019. "GPipe: Efficient Training of Giant Neural Networks Using Pipeline Parallelism." *Advances in Neural Information Processing Systems (NeurIPS)* 32: 103–12.
- Hubara, Itay, Matthieu Courbariaux, Daniel Soudry, Ran El-Yaniv, and Yoshua Bengio. 2018. "Quantized Neural Networks: Training Neural Networks with Low Precision Weights and Activations." *Journal of Machine Learning Research* 18 (187): 1–30.
- Hugging Face. 2026. *Safetensors Documentation*.
- HumanSignal. 2024. *Label Studio: Open Source Data Labeling Platform*.
- Hutter, Frank, Lars Kotthoff, and Joaquin Vanschoren. 2019. *Automated Machine Learning: Methods, Systems, Challenges*. In *Automated Machine Learning*. The Springer Series on Challenges in Machine Learning. Springer International Publishing. <https://doi.org/10.1007/978-3-030-05318-5>.
- Hwu, Wen-mei W. 2011. "Introduction." In *GPU Computing Gems Emerald Edition*. Elsevier. <https://doi.org/10.1016/b978-0-12-384988-5.00064-4>.

- Iandola, Forrest N., Song Han, Matthew W. Moskewicz, Khalid Ashraf, William J. Dally, and Kurt Keutzer. 2016. "SqueezeNet: AlexNet-Level Accuracy with 50x Fewer Parameters and <0.5MB Model Size." *ArXiv Preprint* abs/1602.07360.
- IBM. 2024. *IBM Watson OpenScale: Data Drift Detection*.
- IEEE Standards Association. 2019a. *IEEE 2416-2019: Standard for Power Modeling to Enable System-Level Analysis*.
- IEEE Standards Association. 2019b. *IEEE 754-2019: Standard for Floating-Point Arithmetic*. <https://doi.org/10.1109/IEEESTD.2019.8766229>.
- IEEE Standards Association. 2024. *IEEE Standards Association: Working Groups for AI and ML*.
- Ignatov, Andrey, and Radu Timofte. 2024. *AI Benchmark*.
- InfiniBand Trade Association. 2000. *InfiniBand Architecture Specification Volume 1*. InfiniBand Trade Association.
- Intel Corporation. 2021a. "Intel Advanced Matrix Extensions (Intel AMX)." *Intel Architecture Instruction Set Extensions Programming Reference*.
- Intel Corporation. 2021b. *oneDNN: Intel's Deep Learning Neural Network Library*.
- Intel Corporation. 2026a. *Intel oneAPI Math Kernel Library*.
- Intel Corporation. 2026b. *OpenVINO Documentation*.
- Ioffe, Sergey, and Christian Szegedy. 2015. "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift." *Proceedings of the 32nd International Conference on Machine Learning (ICML)* 37: 448–56.
- ISO. 2024. *ISO/IEC JTC 1/SC 42 Artificial Intelligence*.
- Iterative. 2024. *Data Version Control (DVC)*.
- Jacob, Benoit, Skirmantas Kligys, Bo Chen, Menglong Zhu, Matthew Tang, Andrew Howard, Hartwig Adam, and Dmitry Kalenichenko. 2018. "Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference." *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2704–13. <https://doi.org/10.1109/cvpr.2018.00286>.
- Janapa Reddi, Vijay et al. 2022. "MLPerf Mobile V2. 0: An Industry-Standard Benchmark Suite for Mobile Machine Learning." *Proceedings of Machine Learning and Systems (MLSys)* 4: 806–23.
- Janapa Reddi, Vijay, Brian Plancher, Susan Kennedy, Laurence Moroney, Pete Warden, Lara Suzuki, Anant Agarwal, et al. 2022. "Widening Access to Applied Machine Learning with TinyML." *Harvard Data Science Review* 4 (1). <https://doi.org/10.1162/99608f92.762d171a>.
- Jevons, William Stanley. 1865. *The Coal Question: An Inquiry Concerning the Progress of the Nation, and the Probable Exhaustion of Our Coal Mines*. Macmillan; Co.
- Jia, Xu, Bert De Brabandere, Tinne Tuytelaars, and Luc Van Gool. 2016. "Dynamic Filter Networks." *Advances in Neural Information Processing Systems (NeurIPS)* 29.
- Jia, Yangqing, Evan Shelhamer, Jeff Donahue, Sergey Karayev, Jonathan Long, Ross Girshick, Sergio Guadarrama, and Trevor Darrell. 2014. "Caffe: Convolutional Architecture for Fast Feature Embedding." *Proceedings of the 22nd ACM International Conference on Multimedia*, 675–78. <https://doi.org/10.1145/2647868.2654889>.
- Jia, Zhihao, James Thomas, Todd Warszawski, Mingyu Gao, Matei Zaharia, and Alex Aiken. 2019. "Optimizing DNN Computation with Relaxed Graph Substitutions." *Proceedings of Machine Learning and Systems (MLSys)*.
- Jiang, A. Q., A. Sablayrolles, A. Roux, A. Mensch, B. Savary, C. Bamford, D. S. Chaplot, et al. 2024. "Mixtral of Experts." *arXiv Preprint arXiv:2401.04088*.
- Jiao, Xiaoqi, Yichun Yin, Lifeng Shang, Xin Jiang, Xiao Chen, Linlin Li, Fang Wang, and Qun Liu. 2020. "TinyBERT: Distilling BERT for Natural Language Understanding." *Findings of the Association for Computational Linguistics: EMNLP 2020*, 4163–74. <https://doi.org/10.18653/v1/2020.findings-emnlp.372>.
- Johnson, Jeff, Matthijs Douze, and Hervé Jégou. 2019. "Billion-Scale Similarity Search with GPUs." *IEEE Transactions on Big Data* 7 (3): 535–47. <https://doi.org/10.1109/tbdata.2019.2921572>.
- Johnston, Casey. 2012. "Netflix Never Used Its \$1 Million Algorithm Due to Engineering Costs." *Ars Technica*, April.
- Jordan, T. L. 1982. "A Guide to Parallel Computation and Some Cray-1 Experiences." In *Parallel Computations*. Elsevier. <https://doi.org/10.1016/b978-0-12-592101-5.50006-3>.
- Jouppi, Norman P., Doe Hyun Yoon, Matthew Ashcraft, Mark Gottscho, Thomas B. Jablin, George Kurian, James Laudon, et al. 2021. "Ten Lessons from Three Generations Shaped Google's TPUv4i: Industrial Product." *2021 ACM/IEEE 48th Annual International Symposium on Computer Architecture (ISCA)*, 1–14. <https://doi.org/10.1109/isca52012.2021.00010>.

- Jouppi, Norman P., Cliff Young, Nishant Patil, David Patterson, Gaurav Agrawal, Raminder Bajwa, Sarah Bates, et al. 2017. "In-Datcenter Performance Analysis of a Tensor Processing Unit." *Proceedings of the 44th Annual International Symposium on Computer Architecture, ISCA '17*, 1–12. <https://doi.org/10.1145/3079856.3080246>.
- Jouppi, Norm, George Kurian, Sheng Li, Peter Ma, Rahul Nagarajan, Lifeng Nai, Nishant Patil, et al. 2023. "TPU v4: An Optically Reconfigurable Supercomputer for Machine Learning with Hardware Support for Embeddings." *Proceedings of the 50th Annual International Symposium on Computer Architecture*, 1–14. <https://doi.org/10.1145/3579371.3589350>.
- Jumper, John, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger, Kathryn Tunyasuvunakool, et al. 2021. "Highly Accurate Protein Structure Prediction with AlphaFold." *Nature* 596 (7873): 583–89. <https://doi.org/10.1038/s41586-021-03819-2>.
- Kalamkar, Dhiraj, Dheevatsa Mudigere, Naveen Mellempudi, Dipankar Das, Kunal Banerjee, Sasikanth Avancha, Dharma Teja Vooturi, et al. 2019. *A Study of BFLOAT16 for Deep Learning Training*.
- Kamiran, Faisal, and Toon Calders. 2012. "Data Preprocessing Techniques for Classification Without Discrimination." *Knowledge and Information Systems* 33 (1): 1–33. <https://doi.org/10.1007/s10115-011-0463-8>.
- Kaplan, J., S. McCandlish, T. Henighan, T. B. Brown, B. Chess, R. Child, S. Gray, A. Radford, J. Wu, and D. Amodei. 2020. "Scaling Laws for Neural Language Models." *ArXiv Preprint* abs/2001.08361.
- Karpathy, Andrej. 2017. *Software 2.0*. Medium.
- Kiela, Douwe, Max Bartolo, Yixin Nie, Divyansh Kaushik, Atticus Geiger, Zhengxuan Wu, Bertie Vidgen, et al. 2021. "Dynabench: Rethinking Benchmarking in NLP." *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies* 9: 4110–24. <https://doi.org/10.18653/v1/2021.naacl-main.324>.
- Kilburn, Tom, David B. G. Edwards, Michael J. Lanigan, and Frank H. Sumner. 1962. "One-Level Storage System." *IRE Transactions on Electronic Computers* EC-11 (2): 223–35. <https://doi.org/10.1109/TEC.1962.5219356>.
- Kim, Wonyoung, Meeta S. Gupta, Gu-Yeon Wei, and David Brooks. 2008. "System Level Analysis of Fast, Per-Core DVFS Using on-Chip Switching Regulators." *IEEE International Symposium on High-Performance Computer Architecture (HPCA)*, HPCA '08, 123–34. <https://doi.org/10.1109/hpca.2008.4658633>.
- Kingma, Diederik P., and Jimmy Ba. 2015. "Adam: A Method for Stochastic Optimization." *3rd International Conference on Learning Representations (ICLR)*.
- Kipf, Thomas N., and Max Welling. 2017. "Semi-Supervised Classification with Graph Convolutional Networks." *International Conference on Learning Representations (ICLR)*.
- Kleinberg, Jon, Sendhil Mullainathan, and Manish Raghavan. 2017. "Inherent Trade-Offs in the Fair Determination of Risk Scores." *8th Innovations in Theoretical Computer Science Conference (ITCS 2017)*, Leibniz international proceedings in informatics (LIPIcs), vol. 67: 43:1–23. <https://doi.org/10.4230/LIPIcs.ITCS.2017.43>.
- Kleppmann, Martin. 2016. *Designing Data-Intensive Applications: The Big Ideas Behind Reliable, Scalable, and Maintainable Systems*. O'Reilly Media.
- Kluyver, Thomas, Benjamin Ragan-Kelley, Fernando Pérez, Brian Granger, Matthias Bussonnier, Jonathan Frederic, Kyle Kelley, et al. 2016. "Jupyter Notebooks – a Publishing Format for Reproducible Computational Workflows." In *Positioning and Power in Academic Publishing: Players, Agents and Agendas*, edited by Fernando Loizides and Birgit Schmidt. IOS Press. <https://doi.org/10.3233/978-1-61499-649-1-87>.
- Knight, Will. 2023. *OpenAI's CEO Says the Age of Giant AI Models Is Already over*. WIRED.
- Koh, Pang Wei, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, et al. 2021. "WILDS: A Benchmark of in-the-Wild Distribution Shifts." In *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18–24 July 2021, Virtual Event*, edited by Marina Meila and Tong Zhang, vol. 139. Proceedings of Machine Learning Research. PMLR.
- Koizumi, Yuma, Shoichiro Saito, Hisashi Uematsu, Noboru Harada, and Keisuke Imoto. 2019. "ToyADMOS: A Dataset of Miniature-Machine Operating Sounds for Anomalous Sound Detection." *2019 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, 313–17. <https://doi.org/10.1109/waspaa.2019.8937164>.
- Koomey, Jonathan, Stephen Berard, Marla Sanchez, and Henry Wong. 2011. "Implications of Historical Trends in the Electrical Efficiency of Computing." *IEEE Annals of the History of Computing* 33 (3): 46–54. <https://doi.org/10.1109/mahc.2010.28>.
- Kreuzberger, Dominik, Niklas Kühl, and Sebastian Hirschl. 2023. "Machine Learning Operations (MLOps): Overview, Definition, and Architecture." *IEEE Access* 11: 31866–79. <https://doi.org/10.1109/access.2023.3262138>.
- Krishnamoorthi, Raghuraman. 2018. "Quantizing Deep Convolutional Networks for Efficient Inference: A Whitepaper." *arXiv Preprint arXiv:1806.08342* abs/1806.08342.
- Krishnan, Rayan, Pranav Rajpurkar, and Eric J. Topol. 2022. "Self-Supervised Learning in Medicine and Healthcare." *Nature Biomedical Engineering* 6 (12): 1346–52. <https://doi.org/10.1038/s41551-022-00914-1>.
- Krizhevsky, Alex. 2009. *Learning Multiple Layers of Features from Tiny Images*. University of Toronto.

- Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. 2012. "ImageNet Classification with Deep Convolutional Neural Networks." *Advances in Neural Information Processing Systems (NeurIPS)* 25.
- KServe Community. 2024. *KServe: Highly Scalable and Standards-Based Model Inference Platform on Kubernetes*.
- Kuhn, Max, and Kjell Johnson. 2013. *Applied Predictive Modeling*. Springer New York. <https://doi.org/10.1007/978-1-4614-6849-3>.
- Kullback, Solomon, and Richard A. Leibler. 1951. "On Information and Sufficiency." *The Annals of Mathematical Statistics* 22 (1): 79–86. <https://doi.org/10.1214/aoms/1177729694>.
- Kung, H. T. 1982. "Why Systolic Architectures?" *Computer* 15 (1): 37–46. <https://doi.org/10.1109/mc.1982.1653825>.
- Kung, Hsiang Tsung, and Charles E. Leiserson. 1979. "Systolic Arrays (for VLSI)." *Sparse Matrix Proceedings 1978* 1: 256–82.
- Kuzmin, Andrey, Mart Van Baalen, Yuwei Ren, Markus Nagel, Jorn Peters, and Tijmen Blankevoort. 2022. "FP8 Quantization: The Power of the Exponent." *Advances in Neural Information Processing Systems (NeurIPS)*, 14651–62. <https://doi.org/10.52202/068431-1065>.
- Kuznetsova, Alina, Hassan Rom, Neil Alldrin, Jasper Uijlings, Ivan Krasin, Jordi Pont-Tuset, Shahab Kamali, et al. 2020. "The Open Images Dataset V4: Unified Image Classification, Object Detection, and Visual Relationship Detection at Scale." *International Journal of Computer Vision* 128 (7): 1956–81. <https://doi.org/10.1007/s11263-020-01316-z>.
- Kwon, Woosuk, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. 2023. "Efficient Memory Management for Large Language Model Serving with PagedAttention." *Proceedings of the 29th Symposium on Operating Systems Principles*, 611–26. <https://doi.org/10.1145/3600006.3613165>.
- Kwon, Youngeun, and Minsoo Rhu. 2018. "TensorDIMM: A Practical Near-Memory Processing Architecture for Embeddings and Tensor Operations in Deep Learning." *Proceedings of the 52nd Annual IEEE/ACM International Symposium on Microarchitecture*, 740–53. <https://doi.org/10.1145/3352460.3358284>.
- Labelbox, Inc. 2024. *Labelbox Consensus Documentation*.
- Labs, Grafana. 2024. *Grafana*.
- Lai, Liangzhen, Naveen Suda, and Vikas Chandra. 2018. "CMSIS-NN: Efficient Neural Network Kernels for Arm Cortex-M CPUs." *ArXiv Preprint abs/1801.06601*.
- Lam, Monica D., Edward E. Rothberg, and Michael E. Wolf. 1991. "The Cache Performance and Optimizations of Blocked Algorithms." *Proceedings of the Fourth International Conference on Architectural Support for Programming Languages and Operating Systems - ASPLOS-IV*, 63–74. <https://doi.org/10.1145/106972.106981>.
- Lamport, Leslie, Robert Shostak, and Marshall Pease. 1982. "The Byzantine Generals Problem." *ACM Transactions on Programming Languages and Systems* 4 (3): 382–401. <https://doi.org/10.1145/357172.357176>.
- Lampson, Butler W. 1983. "Hints for Computer System Design." *Proceedings of the Ninth ACM Symposium on Operating Systems Principles*, 33–48. <https://doi.org/10.1145/800217.806614>.
- Landis, J. Richard, and Gary G. Koch. 1977. "The Measurement of Observer Agreement for Categorical Data." *Biometrics* 33 (1): 159–74. <https://doi.org/10.2307/2529310>.
- Lange, Klaus-Dieter. 2009. "Identifying Shades of Green: The SPECpower Benchmarks." *IEEE Computer* 42 (3): 95–97. <https://doi.org/10.1109/mc.2009.84>.
- Lattner, Chris, Mehdi Amini, Uday Bondhugula, Albert Cohen, Andy Davis, Jacques Pienaar, River Riddle, Tatiana Shpeisman, Nicolas Vasilache, and Oleksandr Zinenko. 2020. "MLIR: Scaling Compiler Infrastructure for Domain Specific Computation." *2021 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*, 2–14. <https://doi.org/10.1109/cgo51591.2021.9370308>.
- Lawson, Charles L., Richard J. Hanson, David R. Kincaid, and Fred T. Krogh. 1979. "Basic Linear Algebra Subprograms for Fortran Usage." *ACM Transactions on Mathematical Software* 5 (3): 308–23. <https://doi.org/10.1145/355841.355847>.
- Lazer, David, Ryan Kennedy, Gary King, and Alessandro Vespignani. 2014. "The Parable of Google Flu: Traps in Big Data Analysis." *Science* 343 (6176): 1203–5. <https://doi.org/10.1126/science.1248506>.
- Le Sueur, Etienne, and Gernot Heiser. 2010. "Dynamic Voltage and Frequency Scaling: The Laws of Diminishing Returns." *Proceedings of the 2010 International Conference on Power Aware Computing and Systems*, 1–8.
- Lebedev, Vadim, Yaroslav Ganin, Maksim Rakhuba, Ivan Oseledets, and Victor Lempitsky. 2015. "Speeding up Convolutional Neural Networks Using Fine-Tuned CP-Decomposition." *3rd International Conference on Learning Representations (ICLR)*.
- LeCun, Yann, Yoshua Bengio, and Geoffrey Hinton. 2015. "Deep Learning." *Nature* 521 (7553): 436–44. <https://doi.org/10.1038/nature14539>.
- LeCun, Yann, Bernhard Boser, John S. Denker, Donald Henderson, Richard E. Howard, Wayne Hubbard, and Lawrence D. Jackel. 1989. "Backpropagation Applied to Handwritten Zip Code Recognition." *Neural Computation* 1 (4): 541–51. <https://doi.org/10.1162/neco.1989.1.4.541>.

- LeCun, Yann, Léon Bottou, Yoshua Bengio, and Patrick Haffner. 1998. "Gradient-Based Learning Applied to Document Recognition." *Proceedings of the IEEE* 86 (11): 2278–324. <https://doi.org/10.1109/5.726791>.
- LeCun, Yann, Leon Bottou, Genevieve B. Orr, and Klaus-Robert Müller. 2012. "Efficient BackProp." In *Neural Networks: Tricks of the Trade*, vol. 7700. Lecture Notes in Computer Science. Springer Berlin Heidelberg. https://doi.org/10.1007/978-3-642-35289-8_3.
- LeCun, Yann, John S. Denker, and Sara A. Solla. 1989. "Optimal Brain Damage." *Advances in Neural Information Processing Systems* 2 (NIPS 1989), 598–605.
- Lee, Katherine, Daphne Ippolito, Andrew Nystrom, Chiyuan Zhang, Douglas Eck, Chris Callison-Burch, and Nicholas Carlini. 2022. "Deduplicating Training Data Makes Language Models Better." *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 8424–45. <https://doi.org/10.18653/v1/2022.acl-long.577>.
- Lee, Peter. 2016. *Learning from Tay's Introduction*. The Official Microsoft Blog.
- Lehdonvirta, Vili, Boxi Wu, Zoe Jay Hawkins, Celine Caira, and Lucia Russo. 2025. *Measuring Domestic Public Cloud Compute Availability for Artificial Intelligence*. No. 49. OECD Artificial Intelligence Papers. OECD Publishing. <https://doi.org/10.1787/8602a322-en>.
- Lepikhin, Dmitry, Hyoungho Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. 2021. "GShard: Scaling Giant Models with Conditional Computation and Automatic Sharding." *International Conference on Learning Representations (ICLR)*.
- Leviathan, Yaniv, Matan Kalman, and Yossi Matias. 2023. "Fast Inference from Transformers via Speculative Decoding." *International Conference on Machine Learning (ICML)*, 19274–86.
- Li, Chengwei. 2020. *Estimating the Training Cost of GPT-3*.
- Li, Hao, Asim Kadav, Igor Durdanovic, Hanan Samet, and Hans Peter Graf. 2017. "Pruning Filters for Efficient ConvNets." *International Conference on Learning Representations (ICLR)*.
- Li, Lisha, Kevin G. Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. 2017. "Hyperband: A Novel Bandit-Based Approach to Hyperparameter Optimization." *Journal of Machine Learning Research* 18: 185:1–52.
- Li, Mingzhen, Yi Liu, Xiaoyan Liu, Qingxiao Sun, Xin You, Hailong Yang, Zhongzhi Luan, Lin Gan, Guangwen Yang, and Depei Qian. 2021. "The Deep Learning Compiler: A Comprehensive Survey." *IEEE Transactions on Parallel and Distributed Systems* 32 (3): 708–27. <https://doi.org/10.1109/tpds.2020.3030548>.
- Li, Mu, David G. Andersen, Jun Woo Park, Alexander J. Smola, Amr Ahmed, Vanja Josifovski, James Long, Eugene J. Shekita, and Bor-Ying Su. 2014. "Communication Efficient Distributed Machine Learning with the Parameter Server." *Advances in Neural Information Processing Systems (NeurIPS)* 27: 19–27.
- Li, Ninghui, Tiancheng Li, and Suresh Venkatasubramanian. 2007. "t-Closeness: Privacy Beyond k-Anonymity and l-Diversity." *2007 IEEE 23rd International Conference on Data Engineering*, 106–15. <https://doi.org/10.1109/ICDE.2007.367856>.
- Li, Zhuohan, Lianmin Zheng, Yinmin Zhong, Vincent Liu, Ying Sheng, Xin Jin, Yanping Huang, et al. 2023. "{AlpaServe}: Statistical Multiplexing with Model Parallelism for Deep Learning Serving." *17th USENIX Symposium on Operating Systems Design and Implementation (OSDI 23)*, 663–79.
- Liang, Percy, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, et al. 2022. "Holistic Evaluation of Language Models." *arXiv Preprint arXiv:2211.09110*.
- Lie, Sean. 2021. *Thinking Outside the Die: Architecting the ML Accelerator of the Future*.
- Lighthill, James. 1973. "Artificial Intelligence: A General Survey." In *Artificial Intelligence: A Paper Symposium*. Science Research Council.
- Lin, Ji, Wei-Ming Chen, Yujun Lin, John Cohn, Chuang Gan, and Song Han. 2020. "MCUNet: Tiny Deep Learning on IoT Devices." In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, Virtual*, edited by Hugo Larochelle, Marc'Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin. Curran Associates.
- Lin, Ji, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. 2024. "AWQ: Activation-Aware Weight Quantization for on-Device LLM Compression and Acceleration." *Proceedings of Machine Learning and Systems* 6: 87–100.
- Lin, Tsung-Yi, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C. Lawrence Zitnick. 2014. "Microsoft COCO: Common Objects in Context." In *European Conference on Computer Vision (ECCV)*. Springer. https://doi.org/10.1007/978-3-319-10602-1_48.
- Lindholm, Erik, John Nickolls, Stuart Oberman, and John Montrym. 2008. "NVIDIA Tesla: A Unified Graphics and Computing Architecture." *IEEE Micro* 28 (2): 39–55. <https://doi.org/10.1109/mm.2008.31>.
- Linnainmaa, Seppo. 1970. "The Representation of the Cumulative Rounding Error of an Algorithm as a Taylor Expansion of the Local Rounding Errors." Master's thesis, University of Helsinki.

- Little, John D. C. 1961. "A Proof for the Queuing Formula: $L = \lambda W$." *Operations Research* 9 (3): 383–87. <https://doi.org/10.1287/opre.9.3.383>.
- Liu, Hanxiao, Karen Simonyan, and Yiming Yang. 2019. "DARTS: Differentiable Architecture Search." *International Conference on Learning Representations (ICLR)*.
- Liu, Ze, Yutong Lin, Yue Cao, Han Hu, Yixuan Wei, Zheng Zhang, Stephen Lin, and Baining Guo. 2021. "Swin Transformer: Hierarchical Vision Transformer Using Shifted Windows." *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, 9992–10002. <https://doi.org/10.1109/ICCV48922.2021.00986>.
- Liu, Zhuang, Hanzi Mao, Chao-Yuan Wu, Christoph Feichtenhofer, Trevor Darrell, and Saining Xie. 2022. "A ConvNet for the 2020s." *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 11966–76. <https://doi.org/10.1109/CVPR52688.2022.01167>.
- Loshchilov, Ilya, and Frank Hutter. 2019. "Decoupled Weight Decay Regularization." *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Lowe, David G. 1999. "Object Recognition from Local Scale-Invariant Features." *Proceedings of the Seventh IEEE International Conference on Computer Vision* 2: 1150–1157 vol.2. <https://doi.org/10.1109/iccv.1999.790410>.
- Lu, Yucheng, Shivani Agrawal, Suvinay Subramanian, Oleg Rybakov, Christopher De Sa, and Amir Yazdanbakhsh. 2023. "STEP: Learning n:m Structured Sparsity Masks from Scratch with Precondition." *arXiv Preprint*.
- Lucic, Mario, Karol Kurach, Marcin Michalski, Sylvain Gelly, and Olivier Bousquet. 2018. "Are GANs Created Equal? A Large-Scale Study." *Advances in Neural Information Processing Systems (NeurIPS)* 31.
- Lundberg, Scott M., and Su-In Lee. 2017. "A Unified Approach to Interpreting Model Predictions." *Advances in Neural Information Processing Systems (NeurIPS)* 30: 4765–74.
- Lyons, Richard G. 2011. *Understanding Digital Signal Processing*. 3rd ed. Prentice Hall.
- Machanavajhala, Ashwin, Daniel Kifer, Johannes Gehrke, and Muthu Venkitasubramaniam. 2007. "l-Diversity: Privacy Beyond k-Anonymity." *ACM Transactions on Knowledge Discovery from Data* 1 (3): 3. <https://doi.org/10.1145/1217299.1217302>.
- Mars Climate Orbiter Mishap Investigation Board. 1999. *Mars Climate Orbiter Mishap Investigation Board Phase i Report*. National Aeronautics; Space Administration.
- Maslej, Nestor, Loredana Fattorini, C. Raymond Perrault, Vanessa Parli, Anka Reuel, Erik Brynjolfsson, John Etchemendy, et al. 2024. "Artificial Intelligence Index Report 2024." In CoRR, abs/2405.19522. <https://doi.org/10.48550/ARXIV.2405.19522>.
- Mattson, Peter, Christine Cheng, Cody Coleman, Greg Diamos, Paulius Micekevicius, David Patterson, Hanlin Tang, et al. 2020. "MLPerf Training Benchmark." *Proceedings of Machine Learning and Systems (MLSys)* 2: 336–49.
- Mazumder, Mark. 2021. *TinyML: Multilingual Spoken Words Corpus*. tinyML Talks.
- Mazumder, Mark, Colby Banbury, Xiaozhe Yao, Bojan Karlaš, et al. 2023. "DataPerf: Benchmarks for Data-Centric AI Development." *Advances in Neural Information Processing Systems* 36.
- Mazumder, Mark, Sharad Chitlangia, Colby Banbury, Yiping Kang, Juan Manuel Ciro, Keith Achorn, Daniel Galvez, et al. 2021. "Multilingual Spoken Words Corpus." *Advances in Neural Information Processing Systems 35 (NeurIPS 2021) Datasets and Benchmarks Track*.
- McCarthy, John, Marvin L. Minsky, Nathaniel Rochester, and Claude E. Shannon. 1955. *A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence, August 31, 1955*. No. 4. Vol. 27. Dartmouth College. <https://doi.org/10.1609/aimag.v27i4.1904>.
- McCrory, Dave. 2010. *Data Gravity – in the Clouds*.
- McCulloch, Warren S., and Walter Pitts. 1943. "A Logical Calculus of the Ideas Immanent in Nervous Activity." *The Bulletin of Mathematical Biophysics* 5 (4): 115–33. <https://doi.org/10.1007/bf02478259>.
- McKinsey Global Institute. 2021. *The Internet of Things: Catching up to an Accelerating Opportunity*. McKinsey & Company.
- McMahan, Brendan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera y Arcas. 2017. "Communication-Efficient Learning of Deep Networks from Decentralized Data." *International Conference on Artificial Intelligence and Statistics (AISTATS)*, Proceedings of machine learning research, vol. 54: 1273–82.
- Mellempudi, Naveen, Sudarshan Srinivasan, Dipankar Das, and Bharat Kaul. 2019. "Mixed Precision Training with 8-Bit Floating Point." *arXiv Preprint arXiv:1905.12334*.
- Merity, Stephen. 2016. *The WikiText Long Term Dependency Language Modeling Dataset*.
- Merity, Stephen, Caiming Xiong, James Bradbury, and Richard Socher. 2016. "Pointer Sentinel Mixture Models." *arXiv Preprint arXiv:1609.07843*.

- Merkel, Dirk. 2014. "Docker: Lightweight Linux Containers for Consistent Development and Deployment." *Linux Journal* 2014 (239).
- Michel, Jean-Baptiste, Yuan Kui Shen, Aviva Presser Aiden, Adrian Veres, Matthew K. Gray, William Brockman, Joseph P. Pickett, et al. 2011. "Quantitative Analysis of Culture Using Millions of Digitized Books." *Science* 331 (6014): 176–82. <https://doi.org/10.1126/science.1199644>.
- Michel, Paul, Omer Levy, and Graham Neumann. 2019. "Are Sixteen Heads Really Better Than One?" *Advances in Neural Information Processing Systems (NeurIPS)*.
- Micikevicius, Paulius, Sharan Narang, Jonah Alben, Gregory Diamos, Erich Elsen, David Garcia, Boris Ginsburg, et al. 2017. "Mixed Precision Training." *arXiv Preprint arXiv:1710.03740*.
- Micikevicius, Paulius, Dusan Stolic, Neil Burgess, Marius Cornea, Pradeep Dubey, Richard Grisenthwaite, Sangwon Ha, et al. 2022. "FP8 Formats for Deep Learning." *arXiv Preprint arXiv:2209.05433*.
- Microsoft. 2024a. *Microsoft Azure*.
- Microsoft. 2024b. *ONNX Runtime: Cross-Platform Inference and Training Machine-Learning Accelerator*. GitHub.
- Mikolov, Tomas, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. "Efficient Estimation of Word Representations in Vector Space." *ICLR abs/1301.3781*.
- Mila. 2026. *Open Source Software*.
- Minsky, Marvin, and Seymour A. Papert. 1969. *Perceptrons: An Introduction to Computational Geometry*. The MIT Press.
- Mirhoseini, Azalia, Hieu Pham, Quoc V. Le, Benoit Steiner, Rasmus Larsen, Yuefeng Zhou, Naveen Kumar, Mohammad Norouzi, Samy Bengio, and Jeff Dean. 2017. "Device Placement Optimization with Reinforcement Learning." *International Conference on Machine Learning (ICML)* 70: 2430–39.
- Mitchell, Margaret, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. 2019. "Model Cards for Model Reporting." *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 220–29. <https://doi.org/10.1145/3287560.3287596>.
- Mitchell, Tom M. 1980. *The Need for Biases in Learning Generalizations*. CBM-TR-117. Rutgers University, Department of Computer Science.
- MLCommons. 2024a. *MLPerf Mobile Benchmark*.
- MLCommons. 2024b. *MLPerf Power Measurement*.
- MLCommons. 2024c. *MLPerf Training Benchmark*.
- MLCommons. 2026a. *MLCommons: About Us*.
- MLCommons. 2026b. *MLPerf Client Benchmark*.
- MLflow Project. 2026. *MLflow Model Registry*.
- Moore, G. E. 1998. "Cramming More Components onto Integrated Circuits." *Proceedings of the IEEE* 86 (1): 82–85. <https://doi.org/10.1109/jproc.1998.658762>.
- Moravec, Hans. 1988. *Mind Children: The Future of Robot and Human Intelligence*. Harvard University Press.
- Moritz, Philipp, Robert Nishihara, Stephanie Wang, Alexey Tumanov, Richard Liaw, Eric Liang, Melih Elibol, et al. 2018. "Ray: A Distributed Framework for Emerging AI Applications." *Proceedings of the 13th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 561–77.
- Mosseri, Adam. 2018. *Bringing People Closer Together*. Meta Newsroom.
- Murray, Derek G., Jiří Šimša, Ana Klimovic, and Ihor Indyk. 2021. "Tf.data: A Machine Learning Data Processing Framework." *Proceedings of the VLDB Endowment* 14 (12): 2945–58. <https://doi.org/10.14778/3476311.3476374>.
- Nair, Vinod, and Geoffrey E. Hinton. 2010. "Rectified Linear Units Improve Restricted Boltzmann Machines." *Proceedings of the 27th International Conference on Machine Learning (ICML-10)*, 807–14.
- Nakkiran, Preetum, Gal Kaplun, Yamini Bansal, Tristan Yang, Boaz Barak, and Ilya Sutskever. 2019. "Deep Double Descent: Where Bigger Models and More Data Hurt." *arXiv Preprint arXiv:1912.02292*.
- Narayanan, Deepak, Aaron Harlap, Amar Phanishayee, Vivek Seshadri, Nikhil R. Devanur, Gregory R. Ganger, Phillip B. Gibbons, and Matei Zaharia. 2019. "PipeDream: Generalized Pipeline Parallelism for DNN Training." *Proceedings of the 27th ACM Symposium on Operating Systems Principles*, 1–15. <https://doi.org/10.1145/3341301.3359646>.

- Narayanan, Deepak, Mohammad Shoeybi, Jared Casper, Patrick LeGresley, Mostofa Patwary, Vijay Korthikanti, Dmitri Vainbrand, et al. 2021. "Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM." *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, 1–15. <https://doi.org/10.1145/3458817.3476209>.
- National Aeronautics and Space Administration. 2016. *NASA Systems Engineering Handbook*. NASA/SP-2016-6105 Rev2. National Aeronautics; Space Administration.
- National Transportation Safety Board. 2017. *Collision Between a Car Operating with Automated Vehicle Control Systems and a Tractor-Semitrailer Truck Near Williston, Florida, May 7, 2016*. HAR-17/02. National Transportation Safety Board.
- National Transportation Safety Board. 2019. *Collision Between Vehicle Controlled by Developmental Automated Driving System and Pedestrian, Tempe, Arizona, March 18, 2018*. HAR-19/03. National Transportation Safety Board.
- Naumov, Maxim, Dheevatsa Mudigere, Hao-Jun Michael Shi, Jianyu Huang, Narayanan Sundaraman, Jongsoo Park, Xiaodong Wang, et al. 2019. "Deep Learning Recommendation Model for Personalization and Recommendation Systems." *arXiv Preprint arXiv:1906.00091*.
- Naur, Peter, and Brian Randell, eds. 1969. *Software Engineering: Report of a Conference Sponsored by the NATO Science Committee, Garmisch, Germany, 7–11 October 1968*. Scientific Affairs Division, NATO.
- Nayak, Prateeth, Takuya Higuchi, Anmol Gupta, Shivesh Ranjan, Stephen Shum, Siddharth Sigtia, Erik Marchi, et al. 2022. "Improving Voice Trigger Detection with Metric Learning." *Interspeech 2022*, 1896–900. <https://doi.org/10.21437/interspeech.2022-11160>.
- Netflix. 2024. *Metaflow*.
- Newell, Allen, and Herbert A. Simon. 1976. "Computer Science as Empirical Inquiry: Symbols and Search." *Communications of the ACM* 19 (3): 113–26. <https://doi.org/10.1145/360018.360022>.
- Neyshabur, Behnam, Srinadh Bhojanapalli, David McAllester, and Nati Srebro. 2017. "Exploring Generalization in Deep Learning." *Advances in Neural Information Processing Systems* 30.
- Ng, Andrew. 2021. *MLOps: From Model-Centric to Data-Centric AI*. DeepLearning.AI.
- Nickolls, John, Ian Buck, Michael Garland, and Kevin Skadron. 2008. "Scalable Parallel Programming with CUDA: Is CUDA the Parallel Programming Model That Application Developers Have Been Waiting For?" *Queue* 6 (2): 40–53. <https://doi.org/10.1145/1365490.1365500>.
- Norrie, Thomas, Nishant Patil, Doe Hyun Yoon, George Kurian, Sheng Li, James Laudon, Cliff Young, Norman Jouppi, and David Patterson. 2021. "The Design Process for Google's Training Chips: TPuv2 and TPuv3." *IEEE Micro* 41 (2): 56–63. <https://doi.org/10.1109/mm.2021.3058217>.
- Northcutt, Curtis G., Anish Athalye, and Jonas Mueller. 2021. "Pervasive Label Errors in Test Sets Destabilize Machine Learning Benchmarks." *arXiv Preprint arXiv:2103.14749* 34: 19075–90. <https://doi.org/10.48550/arxiv.2103.14749>.
- NVIDIA. 2017. *Training with Mixed Precision*.
- NVIDIA. 2020a. *Accelerating Matrix Multiplication with Block Sparse Format and NVIDIA Tensor Cores*.
- NVIDIA. 2020b. *TensorFloat-32 in the A100 GPU Accelerates AI Training, HPC up to 20x*.
- NVIDIA. 2024a. *cuBLAS: CUDA Basic Linear Algebra Subprograms*.
- NVIDIA. 2024b. *NVIDIA Omniverse and Simulation*.
- NVIDIA. 2024c. *NVIDIA TensorRT: Programmable Inference Accelerator*.
- NVIDIA. 2024d. *NVIDIA Triton Inference Server*.
- NVIDIA. 2026a. *Lazy Loading: CUDA Programming Guide*.
- NVIDIA. 2026b. *Model Management: NVIDIA Triton Inference Server*.
- NVIDIA. 2026c. *Model Repository: NVIDIA Triton Inference Server*.
- NVIDIA. 2026d. *NVIDIA CUDA Multi-Process Service*.
- NVIDIA. 2026e. *NVIDIA Multi-Instance GPU User Guide*.
- NVIDIA. 2026f. *NVIDIA Quantum-X800 InfiniBand Platform*. NVIDIA product documentation.
- NVIDIA. 2026g. *NVIDIA TensorRT-LLM*.
- NVIDIA Corporation. 2017. *NVIDIA Tesla V100 GPU Architecture*. NVIDIA Whitepaper.

- NVIDIA Corporation. 2018. *NVIDIA Tesla T4 Tensor Core GPU*. NVIDIA product documentation.
- NVIDIA Corporation. 2020a. *NVIDIA A100 Tensor Core GPU Architecture*. NVIDIA Whitepaper, V1.0.
- NVIDIA Corporation. 2020b. *NVIDIA V100 Tensor Core GPU Datasheet*. NVIDIA product documentation.
- NVIDIA Corporation. 2020c. "NVLink: Scalable High-Performance Interconnect." *NVIDIA Technical Report* 2020.
- NVIDIA Corporation. 2021a. *NVIDIA A100 Tensor Core GPU Datasheet*. NVIDIA product documentation.
- NVIDIA Corporation. 2021b. *NVIDIA cuDNN Developer Guide*.
- NVIDIA Corporation. 2023. *NVIDIA cuDNN Developer Guide, Version 8.9.6*.
- NVIDIA Corporation. 2024. *NVIDIA Blackwell Architecture*. NVIDIA product documentation.
- NVIDIA Corporation. 2026a. *How TensorRT Works*.
- NVIDIA Corporation. 2026b. *Working with Dynamic Shapes*.
- Oakden-Rayner, Luke, Jared Dunnmon, Gustavo Carneiro, and Christopher Re. 2020. "Hidden Stratification Causes Clinically Meaningful Failures in Machine Learning for Medical Imaging." *Proceedings of the ACM Conference on Health, Inference, and Learning*, 151–59. <https://doi.org/10.1145/3368555.3384468>.
- Obermeyer, Ziad, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. 2019. "Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations." *Science* 366 (6464): 447–53. <https://doi.org/10.1126/science.aax2342>.
- Olston, Christopher, Noah Fiedel, Kiril Gorovoy, Jeremiah Harmsen, Li Lao, Fangwei Li, Vinu Rajashekhar, Sukriti Ramesh, and Jordan Soyke. 2017. "TensorFlow-Serving: Flexible, High-Performance ML Serving." *CoRR* abs/1712.06139. <https://doi.org/10.48550/arXiv.1712.06139>.
- ONNX Contributors. 2019. *ONNX: Open Neural Network Exchange*.
- OpenAI. 2023a. *GPT-4*.
- OpenAI. 2023b. *Introducing ChatGPT and Whisper APIs*.
- OpenAI. 2023c. *New Models and Developer Products Announced at DevDay*.
- OpenAI. 2024. *GPT-4o Mini: Advancing Cost-Efficient Intelligence*.
- OpenAI Developer Community. 2024. *Announcing GPT-4o in the API*.
- OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, et al. 2023. "GPT-4 Technical Report." *arXiv Preprint arXiv:2303.08774*, ahead of print. <https://doi.org/10.48550/arXiv.2303.08774>.
- OpenBLAS Project. 2026. *OpenBLAS: An Optimized BLAS Library*.
- Orr, Laurel, Atindriyo Sanyal, Xiao Ling, Karan Goel, and Megan Leszczynski. 2021. "Managing ML Pipelines: Feature Stores and the Coming Wave of Embedding Ecosystems." *Proceedings of the VLDB Endowment* 14 (12): 3178–81. <https://doi.org/10.14778/3476311.3476402>.
- Owens, John D., Mike Houston, David Luebke, Simon Green, John E. Stone, and James C. Phillips. 2008. "GPU Computing." *Proceedings of the IEEE* 96 (5): 879–99. <https://doi.org/10.1109/jproc.2008.917757>.
- Paley, Andrei, Raoul-Gabriel Urma, and Neil D. Lawrence. 2022. "Challenges in Deploying Machine Learning: A Survey of Case Studies." *ACM Computing Surveys* 55 (6): 1–29. <https://doi.org/10.1145/3533378>.
- Palmer, John F. 1980. "The Intel 8087 Numeric Data Processor." *Proceedings of the 1980 National Computer Conference (AFIPS)*, 887–93. <https://doi.org/10.1145/1500518.1500674>.
- Papermill Project. 2026. *Papermill Documentation*.
- Papineni, Kishore, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. "Bleu: A Method for Automatic Evaluation of Machine Translation." *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics - ACL '02*, 311–18. <https://doi.org/10.3115/1073083.1073135>.
- Pareto, Vilfredo. 1896. *Cours d'économie Politique*. F. Rouge.
- Park, Daniel S., William Chan, Yu Zhang, Chung-Cheng Chiu, Barret Zoph, Ekin D. Cubuk, and Quoc V. Le. 2019. "SpecAugment: A Simple Data Augmentation Method for Automatic Speech Recognition." *Interspeech 2019*, 2613–17. <https://doi.org/10.21437/interspeech.2019-2680>.
- Paszke, Adam, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, et al. 2019. "PyTorch: An Imperative Style, High-Performance Deep Learning Library." *Advances in Neural Information Processing Systems (NeurIPS)* 32: 8024–35.

- Patel, Dylan, and Gerald Wong. 2023. *GPT-4 Architecture, Infrastructure, Training Dataset, Costs, Vision, MoE*. SemiAnalysis Blog.
- Patterson, David, Joseph Gonzalez, Quoc Le, Chen Liang, Lluís-Miquel Munguia, Daniel Rothchild, David So, Maud Texier, and Jeff Dean. 2021. "Carbon Emissions and Large Neural Network Training." *arXiv Preprint arXiv:2104.10350*.
- Paul, Mansheej, Surya Ganguli, and Gintare Karolina Dziugaite. 2021. "Deep Learning on a Data Diet: Finding Important Examples Early in Training." *Advances in Neural Information Processing Systems (NeurIPS)* 34: 20596–607.
- Pham, Hieu, Melody Y. Guan, Barret Zoph, Quoc V. Le, and Jeff Dean. 2018. "Efficient Neural Architecture Search via Parameter Sharing." *Proceedings of the 35th International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 80: 4095–104.
- Pichai, Sundar. 2023. *Introducing Gemini: Our Largest and Most Capable AI Model*.
- Pineau, Joelle, Philippe Vincent-Lamarre, Koustuv Sinha, Vincent Larivière, Alina Beygelzimer, Florence d'Alché-Buc, Emily Fox, and Hugo Larochelle. 2021. "Improving Reproducibility in Machine Learning Research (a Report from the NeurIPS 2019 Reproducibility Program)." *Journal of Machine Learning Research* 22 (164): 1–20.
- Pope, Reiner, Sholto Douglas, Aakanksha Chowdhery, Jacob Devlin, James Bradbury, Jonathan Heek, Kefan Xiao, Shivani Agrawal, and Jeff Dean. 2023. "Efficiently Scaling Transformer Inference." *Proceedings of Machine Learning and Systems (MLSys)* 5: 606–24.
- Prefect Technologies, Inc. 2024. *Prefect: Workflow Orchestration Framework for Python*.
- Project Jupyter. 2024. *Project Jupyter*.
- Project Jupyter. 2026. *nbconvert: Convert Notebooks to Other Formats*.
- Protocol Buffers Authors. 2026. *Overview: Protocol Buffers*.
- Pushkarna, Mahima, Andrew Zaldivar, and Oddur Kjartansson. 2022. "Data Cards: Purposeful and Transparent Dataset Documentation for Responsible AI." *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, 1776–826. <https://doi.org/10.1145/3531146.3533231>.
- Putnam, Andrew, Adrian M. Caulfield, Eric S. Chung, Derek Chiou, Kypros Constantinides, John Demme, Hadi Esmaeilzadeh, et al. 2014. "A Reconfigurable Fabric for Accelerating Large-Scale Datacenter Services." *ACM SIGARCH Computer Architecture News* 42 (3): 13–24. <https://doi.org/10.1145/2678373.2665678>.
- PyTorch. 2022. *Accelerating Neural Network Training with Sparse Tensors*.
- PyTorch Contributors. 2026a. *PyTorch Autograd Mechanics*.
- PyTorch Contributors. 2026b. *Tensor Views*.
- PyTorch Contributors. 2026c. *torch.export*.
- PyTorch Contributors. 2026d. *TorchScript*.
- Qi, Chen, Shibo Shen, Rongpeng Li, Zhifeng Zhao, Qing Liu, Jing Liang, and Honggang Zhang. 2021. "An Efficient Pruning Scheme of Deep Neural Networks for Internet of Things Applications." *EURASIP Journal on Advances in Signal Processing* 2021 (1): 31. <https://doi.org/10.1186/s13634-021-00744-4>.
- Quiñero-Candela, Joaquin, Masashi Sugiyama, Anton Schwaighofer, and Neil D. Lawrence, eds. 2009. *Dataset Shift in Machine Learning*. Neural Information Processing Series. The MIT Press.
- Rachwan, John, Daniel Zügner, Bertrand Charpentier, Simon Geisler, Morgane Ayle, and Stephan Günnemann. 2022. "Winning the Lottery Ahead of Time: Efficient Early Network Pruning." *International Conference on Machine Learning (ICML)*, 18293–309.
- Radford, Alec, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, et al. 2021. "Learning Transferable Visual Models from Natural Language Supervision." *International Conference on Machine Learning (ICML)*, 8748–63.
- Radford, Alec, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. *Language Models Are Unsupervised Multitask Learners*. OpenAI.
- Raffel, C., N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu. 2020. "Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer." *Journal of Machine Learning Research* 21 (140): 1–67.
- Raichle, Marcus E., and Debra A. Gusnard. 2002. "Appraising the Brain's Energy Budget." *Proceedings of the National Academy of Sciences* 99 (16): 10237–39. <https://doi.org/10.1073/pnas.172399499>.
- Rainio, Oona, Jarmo Teuho, and Riku Klén. 2024. "Evaluation Metrics and Statistical Tests for Machine Learning." *Scientific Reports* 14 (1): 6086. <https://doi.org/10.1038/s41598-024-56706-x>.

- Raji, Inioluwa Deborah, and Joy Buolamwini. 2019. "Actionable Auditing: Investigating the Impact of Publicly Naming Biased Performance Results of Commercial AI Products." *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society*, 429–35. <https://doi.org/10.1145/3306618.3314244>.
- Rajpurkar, Pranav, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. "SQuAD: 100,000+ Questions for Machine Comprehension of Text." *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, 2383–92. <https://doi.org/10.18653/v1/d16-1264>.
- Ranjit, Mercy Prasanna, Gopinath Ganapathy, Kalaivani Sridhar, and Vikram Arumugham. 2019. "Efficient Deep Learning Hyperparameter Tuning Using Cloud Infrastructure: Intelligent Distributed Hyperparameter Tuning with Bayesian Optimization in the Cloud." *2019 IEEE 12th International Conference on Cloud Computing (CLOUD)*, 520–22. <https://doi.org/10.1109/cloud.2019.00097>.
- Rastegari, Mohammad, Vicente Ordonez, Joseph Redmon, and Ali Farhadi. 2016. "XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks." In *European Conference on Computer Vision (ECCV)*. Springer. https://doi.org/10.1007/978-3-319-46493-0_32.
- Ratner, A., S. H. Bach, H. Ehrenberg, J. Fries, S. Wu, and C. Ré. 2017. "Snorkel: Rapid Training Data Creation with Weak Supervision." *Proceedings of the VLDB Endowment* 11 (3): 269–82. <https://doi.org/10.14778/3157794.3157797>.
- Ratner, Alex, Braden Hancock, Jared Dunnmon, Roger Goldman, and Christopher Ré. 2018. "Snorkel MeTaL: Weak Supervision for Multi-Task Learning." *Proceedings of the Second Workshop on Data Management for End-to-End Machine Learning*, 1–4. <https://doi.org/10.1145/3209889.3209898>.
- Ratner, Alex, Paroma Varma, Braden Hancock, and Christopher Ré. 2019. *Weak Supervision: A New Programming Paradigm for Machine Learning*. Stanford AI Lab Blog.
- Reagen, Brandon, Robert Adolf, Paul Whatmough, Gu-Yeon Wei, and David Brooks. 2017. *Deep Learning for Computer Architects. Synthesis Lectures on Computer Architecture*. Springer International Publishing. <https://doi.org/10.1007/978-3-031-01756-8>.
- Real, Esteban, Alok Aggarwal, Yanping Huang, and Quoc V. Le. 2019. "Regularized Evolution for Image Classifier Architecture Search." *Proceedings of the AAAI Conference on Artificial Intelligence* 33 (01): 4780–89. <https://doi.org/10.1609/aaai.v33i01.33014780>.
- Recht, Benjamin, Rebecca Roelofs, Ludwig Schmidt, and Vaishal Shankar. 2019. "Do ImageNet Classifiers Generalize to ImageNet?" *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 5389–400.
- Reddi, Vijay Janapa, Christine Cheng, David Kanter, Peter Mattson, Guenther Schmuelling, Carole-Jean Wu, Brian Anderson, et al. 2019. "MLPerf Inference Benchmark." *2020 ACM/IEEE 47th Annual International Symposium on Computer Architecture (ISCA)*, 446–59. <https://doi.org/10.1109/isca45697.2020.00045>.
- Ren, Pengzhen, Yun Xiao, Xiaojun Chang, Po-Yao Huang, Zhihui Li, Brij B. Gupta, Xiaojiang Chen, and Xin Wang. 2021. "A Survey of Deep Active Learning." *ACM Computing Surveys* 54 (9): 1–40. <https://doi.org/10.1145/3472291>.
- Ribeiro, M. H., R. Ottoni, R. West, V. A. F. Almeida, and Jr. Meira Wagner. 2020. "Auditing Radicalization Pathways on YouTube." *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 131–41. <https://doi.org/10.1145/3351095.3372879>.
- Ribeiro, Marco Tulio, Sameer Singh, and Carlos Guestrin. 2016. "Why Should I Trust You?: Explaining the Predictions of Any Classifier." *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1135–44. <https://doi.org/10.1145/2939672.2939778>.
- Ribeiro, Marco Tulio, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. "Beyond Accuracy: Behavioral Testing of NLP Models with CheckList." *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 4902–12. <https://doi.org/10.18653/v1/2020.acl-main.442>.
- Roesch, Jared, Steven Lyubomirsky, Logan Weber, Josh Pollock, Marisa Kirisame, Tianqi Chen, and Zachary Tatlock. 2018. "Relay: A New IR for Machine Learning Frameworks." *Proceedings of the 2nd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages*, 58–68. <https://doi.org/10.1145/3211346.3211348>.
- Rombach, Robin, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Bjorn Ommer. 2022. "High-Resolution Image Synthesis with Latent Diffusion Models." *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 10674–85. <https://doi.org/10.1109/cvpr52688.2022.01042>.
- Romero, Francisco, Qian Li, Neeraja J. Yadwadkar, and Christos Kozyrakis. 2021. "INFaaS: Automated Model-Less Inference Serving." *2021 USENIX Annual Technical Conference (USENIX ATC 21)*, 397–411.
- Rosenblatt, Frank. 1957. *The Perceptron: A Perceiving and Recognizing Automaton*. Report Nos. 85-460-1. Cornell Aeronautical Laboratory.
- Rosenblatt, Frank. 1958. "The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain." *Psychological Review* 65 (6): 386–408. <https://doi.org/10.1037/h0042519>.
- Royce, Winston W. 1970. "Managing the Development of Large Software Systems." *Proceedings of IEEE WESCON* 26: 328–88.

- Rumelhart, David E., Geoffrey E. Hinton, and Ronald J. Williams. 1986. "Learning Representations by Back-Propagating Errors." *Nature* 323 (6088): 533–36. <https://doi.org/10.1038/323533a0>.
- Russakovsky, Olga, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, et al. 2015. "ImageNet Large Scale Visual Recognition Challenge." *International Journal of Computer Vision* 115 (3): 211–52. <https://doi.org/10.1007/s11263-015-0816-y>.
- Sabour, Sara, Nicholas Frosst, and Geoffrey E Hinton. 2017. "Dynamic Routing Between Capsules." *Advances in Neural Information Processing Systems (NeurIPS)* 30.
- Sambasivan, Nithya, Shivani Kapania, Hannah Highfill, Diana Akrong, Praveen Paritosh, and Lora M Aroyo. 2021. "'Everyone Wants to Do the Model Work, Not the Data Work': Data Cascades in High-Stakes AI." *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, 1–15. <https://doi.org/10.1145/3411764.3445518>.
- Samuel, A. L. 1959. "Some Studies in Machine Learning Using the Game of Checkers." *IBM Journal of Research and Development* 3 (3): 210–29. <https://doi.org/10.1147/rd.33.0210>.
- Sandler, Mark, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, and Liang-Chieh Chen. 2018. "MobileNetV2: Inverted Residuals and Linear Bottlenecks." *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 4510–20. <https://doi.org/10.1109/cvpr.2018.00474>.
- Sanh, Victor, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2019. "DistilBERT, a Distilled Version of BERT: Smaller, Faster, Cheaper and Lighter." *arXiv Preprint arXiv:1910.01108*, ahead of print. <https://doi.org/10.48550/arXiv.1910.01108>.
- Savage, Sam L. 2009. *The Flaw of Averages: Why We Underestimate Risk in the Face of Uncertainty*. John Wiley & Sons.
- Scale AI, Inc. 2024. *Scale AI Advanced Settings*.
- Schelter, Sebastian, Matthias Boehm, Johannes Kirschnick, Kostas Tzoumas, and Gunnar Ratsch. 2018. "Automating Large-Scale Machine Learning Model Management." *Proceedings of the 2018 IEEE International Conference on Data Engineering (ICDE)*, 137–48.
- Schwartz, Roy, Jesse Dodge, Noah A. Smith, and Oren Etzioni. 2020. "Green AI." *Communications of the ACM* 63 (12): 54–63. <https://doi.org/10.1145/3381831>.
- scikit-learn developers. 2024a. *Feature Selection — Scikit-Learn Documentation*. Scikit-learn Documentation.
- scikit-learn developers. 2024b. *Sklearn.metrics.confusion_matrix — Scikit-Learn Documentation*. Scikit-learn Documentation.
- Scott, Colin. 2012. "Numbers Every Programmer Should Know by Year."
- Sculley, D., Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-François Crespo, and Dan Dennison. 2015. "Hidden Technical Debt in Machine Learning Systems." *Advances in Neural Information Processing Systems (NeurIPS)* 28: 2503–11.
- Sener, Ozan, and Silvio Savarese. 2018. "Active Learning for Convolutional Neural Networks: A Core-Set Approach." *International Conference on Learning Representations (ICLR)*.
- Sennrich, Rico, Barry Haddow, and Alexandra Birch. 2016. "Improving Neural Machine Translation Models with Monolingual Data." *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 86–96. <https://doi.org/10.18653/v1/P16-1009>.
- Sergeev, Alexander, and Mike Del Balso. 2018. "Horovod: Fast and Easy Distributed Deep Learning in TensorFlow." *CoRR* abs/1802.05799.
- Settles, Burr. 2009. *Active Learning Literature Survey*. Computer Sciences Technical Report 1648. University of Wisconsin-Madison.
- Settles, Burr. 2012. *Active Learning*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Springer Cham. <https://doi.org/10.1007/978-3-031-01560-1>.
- Sevilla, Jaime, Lennart Heim, Anson Ho, Tamay Besiroglu, Marius Hobbhahn, and Pablo Villalobos. 2022. "Compute Trends Across Three Eras of Machine Learning." *2022 International Joint Conference on Neural Networks (IJCNN)*, 1–8. <https://doi.org/10.1109/ijcnn55064.2022.9891914>.
- Shah, Jay, Ganesh Bikshandi, Ying Zhang, Vijay Thakkar, Pradeep Ramani, and Tri Dao. 2024. "FlashAttention-3: Fast and Accurate Attention with Asynchrony and Low-Precision." *Advances in Neural Information Processing Systems* 37 (NeurIPS), 68658–85. <https://doi.org/10.52202/079017-2193>.
- Shannon, Claude E. 1948. "A Mathematical Theory of Communication." *Bell System Technical Journal* 27 (3): 379–423. <https://doi.org/10.1002/j.1538-7305.1948.tb01338.x>.
- Shazeer, Noam, Azalia Mirhoseini, Krzysztof Maziarsz, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. 2017. "Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer." *International Conference on Learning Representations (ICLR)*.

- Shazeer, Noam, and Mitchell Stern. 2018. "Adafactor: Adaptive Learning Rates with Sublinear Memory Cost." *International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 80: 4596–604.
- Shen, Haichen, Lequn Chen, Yuchen Jin, Liangyu Zhao, Bingyu Kong, Matthai Philipose, Arvind Krishnamurthy, and Ravi Sundaram. 2019. "Nexus: A GPU Cluster Engine for Accelerating DNN-Based Video Analysis." *Proceedings of the 27th ACM Symposium on Operating Systems Principles*, 322–37. <https://doi.org/10.1145/3341301.3359658>.
- Shen, Sheng, Zhen Dong, Jiayu Ye, Linjian Ma, Zhewei Yao, Amir Gholami, Michael W. Mahoney, and Kurt Keutzer. 2020. "Q-BERT: Hessian Based Ultra Low Precision Quantization of BERT." *Proceedings of the AAAI Conference on Artificial Intelligence* 34 (05): 8815–21. <https://doi.org/10.1609/aaai.v34i05.6409>.
- Sheng, Victor S., and Jing Zhang. 2019. "Machine Learning with Crowdsourcing: A Brief Summary of the Past Research and Future Directions." *Proceedings of the AAAI Conference on Artificial Intelligence* 33 (01): 9837–43. <https://doi.org/10.1609/aaai.v33i01.33019837>.
- Shi, Weisong, Jie Cao, Quan Zhang, Youhuizi Li, and Lanyu Xu. 2016. "Edge Computing: Vision and Challenges." *IEEE Internet of Things Journal* 3 (5): 637–46. <https://doi.org/10.1109/jiot.2016.2579198>.
- Shneiderman, B. 2022. *Human-Centered AI*. Oxford University Press. <https://doi.org/10.1093/oso/9780192845290.001.0001>.
- Shoeybi, Mohammad, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. 2019. "Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism." *arXiv Preprint arXiv:1909.08053*.
- Shokri, Reza, Marco Stronati, Congzheng Song, and Vitaly Shmatikov. 2017. "Membership Inference Attacks Against Machine Learning Models." *2017 IEEE Symposium on Security and Privacy (SP)*, 3–18. <https://doi.org/10.1109/SP.2017.41>.
- Shorten, Connor, and Taghi M. Khoshgoftaar. 2019. "A Survey on Image Data Augmentation for Deep Learning." *Journal of Big Data* 6 (1): 1–48. <https://doi.org/10.1186/s40537-019-0197-0>.
- Shortliffe, Edward H., Randall Davis, Stanton G. Axline, Bruce G. Buchanan, C.Cordell Green, and Stanley N. Cohen. 1975. "Computer-Based Consultations in Clinical Therapeutics: Explanation and Rule Acquisition Capabilities of the MYCIN System." *Computers and Biomedical Research* 8 (4): 303–20. [https://doi.org/10.1016/0010-4809\(75\)90009-9](https://doi.org/10.1016/0010-4809(75)90009-9).
- Shumailov, Ilia, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. 2024. "AI Models Collapse When Trained on Recursively Generated Data." *Nature* 631 (8022): 755–59. <https://doi.org/10.1038/s41586-024-07566-y>.
- Sifre, Laurent, and Stéphane Mallat. 2014. "Rigid-Motion Scattering for Image Classification." *arXiv Preprint arXiv:1403.1687*, ahead of print. <https://doi.org/10.48550/arXiv.1403.1687>.
- Silver, David, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, et al. 2016. "Mastering the Game of Go with Deep Neural Networks and Tree Search." *Nature* 529 (7587): 484–89. <https://doi.org/10.1038/nature16961>.
- Silver, David, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, et al. 2017. "Mastering the Game of Go Without Human Knowledge." *Nature* 550 (7676): 354–59. <https://doi.org/10.1038/nature24270>.
- Simonyan, Karen, and Andrew Zisserman. 2015. "Very Deep Convolutional Networks for Large-Scale Image Recognition." *arXiv Preprint*.
- Smith, Steven W. 1997. "Digital Signal Processing." In *Digital Signal Processing Demystified*. Elsevier. <https://doi.org/10.1016/b978-187870716-1/50004-4>.
- Snow, Rion, Brendan O'Connor, Daniel Jurafsky, and Andrew Y. Ng. 2008. "Cheap and Fast—but Is It Good? Evaluating Non-Expert Annotations for Natural Language Tasks." *Proceedings of the Conference on Empirical Methods in Natural Language Processing - EMNLP '08*, 254. <https://doi.org/10.3115/1613715.1613751>.
- Sohn, Kihyuk, David Berthelot, Nicholas Carlini, Zizhao Zhang, Han Zhang, Colin A Raffel, Ekin Dogus Cubuk, Alexey Kurakin, and Chun-Liang Li. 2020. "FixMatch: Simplifying Semi-Supervised Learning with Consistency and Confidence." *Advances in Neural Information Processing Systems (NeurIPS)* 33: 596–608.
- Sokolova, Marina, and Guy Lapalme. 2009. "A Systematic Analysis of Performance Measures for Classification Tasks." *Information Processing & Management* 45 (4): 427–37. <https://doi.org/10.1016/j.ipm.2009.03.002>.
- Sorscher, Ben, Robert Geirhos, Shashank Shekhar, Surya Ganguli, and Ari S. Morcos. 2022. "Beyond Neural Scaling Laws: Beating Power Law Scaling via Data Pruning." *Advances in Neural Information Processing Systems (NeurIPS)* 35: 19523–36. <https://doi.org/10.52202/068431-1419>.
- Soviany, Petru, Radu Tudor Ionescu, Paolo Rota, and Nicu Sebe. 2022. "Curriculum Learning: A Survey." *International Journal of Computer Vision* 130 (6): 1526–65. <https://doi.org/10.1007/s11263-022-01611-x>.
- Srivastava, Nitish, Geoffrey E. Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. 2014. "Dropout: A Simple Way to Prevent Neural Networks from Overfitting." *Journal of Machine Learning Research* 15 (1): 1929–58.
- Statista Research Department. 2024. *Number of Sent and Received e-Mails Per Day Worldwide from 2017 to 2027*. Statista.

- Stephens, Nigel, Stuart Biles, Matthias Boettcher, Jacob Eapen, Mbou Eyole, Giacomo Gabrielli, Matt Horsnell, et al. 2017. "The ARM Scalable Vector Extension." *IEEE Micro* 37 (2): 26–39. <https://doi.org/10.1109/mm.2017.35>.
- Stonebraker, Mike, Daniel J. Abadi, Adam Batkin, Xuedong Chen, Mitch Cherniack, Miguel Ferreira, Edmond Lau, et al. 2018. "C-Store: A Column-Oriented DBMS." In *Making Databases Work: The Pragmatic Wisdom of Michael Stonebraker*. Association for Computing Machinery. <https://doi.org/10.1145/3226595.3226638>.
- Strassen, Volker. 1969. "Gaussian Elimination Is Not Optimal." *Numerische Mathematik* 13 (4): 354–56. <https://doi.org/10.1007/bf02165411>.
- Strathern, Marilyn. 1997. "'Improving Ratings': Audit in the British University System." *European Review* 5 (3): 305–21. [https://doi.org/10.1002/\(SICI\)1234-981X\(199707\)5:3<305::AID-EURO184>3.0.CO;2-4](https://doi.org/10.1002/(SICI)1234-981X(199707)5:3<305::AID-EURO184>3.0.CO;2-4).
- Stromberg, Joseph. 2015. *When Was the First Traffic Light Installed? Today in 1914*. Vox.
- Strubell, Emma, Ananya Ganesh, and Andrew McCallum. 2019. "Energy and Policy Considerations for Deep Learning in NLP." *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, 3645–50. <https://doi.org/10.18653/v1/p19-1355>.
- Sullivan, Gary J., Jens-Rainer Ohm, Woo-Jin Han, and Thomas Wiegand. 2012. "Overview of the High Efficiency Video Coding (HEVC) Standard." *IEEE Transactions on Circuits and Systems for Video Technology* 22 (12): 1649–68. <https://doi.org/10.1109/tcsvt.2012.2221191>.
- Sun, Pei, Henrik Kretschmar, Xerxes Dotiwalla, Aurelien Chouard, Vijaysai Patnaik, Paul Tsui, James Guo, et al. 2020. "Scalability in Perception for Autonomous Driving: Waymo Open Dataset." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2446–54. <https://doi.org/10.1109/CVPR42600.2020.00252>.
- Supreme Court of the United States. 1971. *Griggs v. Duke Power Co.*, 401 U.S. 424. Legal decision.
- Sutskever, Ilya, Oriol Vinyals, and Quoc V. Le. 2014. "Sequence to Sequence Learning with Neural Networks." *Advances in Neural Information Processing Systems (NeurIPS)* 27: 3104–12.
- Sutton, Richard S. 2019. "The Bitter Lesson." *Incompleteideas.net* 43.
- Sweeney, Latanya. 2002. "K-ANONYMITY: A MODEL for PROTECTING PRIVACY." *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems* 10 (05): 557–70. <https://doi.org/10.1142/s0218488502001648>.
- Systems, Cerebras. 2021. "Wafer-Scale Deep Learning Acceleration with the Cerebras CS-2." *Cerebras Technical Paper* 2021.
- Sze, Vivienne, Yu-Hsin Chen, Tien-Ju Yang, and Joel S. Emer. 2017. "Efficient Processing of Deep Neural Networks: A Tutorial and Survey." *Proceedings of the IEEE* 105 (12): 2295–329. <https://doi.org/10.1109/jproc.2017.2761740>.
- Szegedy, Christian, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. 2015. "Going Deeper with Convolutions." *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 1–9. <https://doi.org/10.1109/cvpr.2015.7298594>.
- Tan, Mingxing, Bo Chen, Ruoming Pang, Vijay Vasudevan, Mark Sandler, Andrew Howard, and Quoc V. Le. 2019. "MnasNet: Platform-Aware Neural Architecture Search for Mobile." *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2815–23. <https://doi.org/10.1109/cvpr.2019.00293>.
- Tan, Mingxing, and Quoc V. Le. 2019. "EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks." *International Conference on Machine Learning (ICML)*, 6105–14.
- Team, The Theano Development, Rami Al-Rfou, Guillaume Alain, Amjad Almahairi, Christof Angermueller, Dzmitry Bahdanau, Nicolas Ballas, et al. 2016. "Theano: A Python Framework for Fast Computation of Mathematical Expressions." *arXiv Preprint*.
- Teerapittayanon, Surat, Bradley McDanel, and H. T. Kung. 2017. "BranchyNet: Fast Inference via Early Exiting from Deep Neural Networks." *2016 23rd International Conference on Pattern Recognition (ICPR)*, 2464–69. <https://doi.org/10.1109/icpr.2016.7900006>.
- Telgarsky, Matus. 2016. "Benefits of Depth in Neural Networks." *Proceedings of the 29th Annual Conference on Learning Theory*, 1517–39.
- TensorFlow Developers. 2024a. *Better Performance with tf.function*.
- TensorFlow Developers. 2024b. *Optimize TensorFlow Performance Using the Profiler*.
- Tesla, Inc. 2021. *Tesla Dojo Technology: A Guide to Tesla's Configurable Floating Point Formats & Arithmetic*. Tesla AI Day, Whitepaper.
- Tesler, Larry. 1984. *The Law of Conservation of Complexity*. Web page.
- Thyagarajan, Aditya, Eliás Snorrason, Curtis G. Northcutt, and Jonas Mueller. 2022. "Identifying Incorrect Annotations in Multi-Label Classification Data." *CoRR abs/2211.13895*. <https://doi.org/10.48550/ARXIV.2211.13895>.

- Tibshirani, Robert. 1996. "Regression Shrinkage and Selection via the Lasso." *Journal of the Royal Statistical Society: Series B (Methodological)* 58 (1): 267–88. <https://doi.org/10.1111/j.2517-6161.1996.tb02080.x>.
- Tillet, Philippe, H. T. Kung, and David Cox. 2019. "Triton: An Intermediate Language and Compiler for Tiled Neural Network Computations." *Proceedings of the 3rd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages*, 10–19. <https://doi.org/10.1145/3315508.3329973>.
- Toneva, Mariya, Alessandro Sordoni, Remi Tachet des Combes, Adam Trischler, Yoshua Bengio, and Geoffrey J Gordon. 2019. "An Empirical Study of Example Forgetting During Deep Neural Network Learning." *International Conference on Learning Representations (ICLR)*.
- Torvalds, Linus, and Junio Hamano. 2024. *Git*.
- Touvron, Hugo, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, et al. 2023. "LLaMA: Open and Efficient Foundation Language Models." *arXiv Preprint arXiv:2302.13971*.
- Touvron, Hugo, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, et al. 2023. "Llama 2: Open Foundation and Fine-Tuned Chat Models." *arXiv Preprint arXiv:2307.09288*.
- Tramèr, Florian, Fan Zhang, Ari Juels, Michael K Reiter, and Thomas Ristenpart. 2016. "Stealing Machine Learning Models via Prediction APIs." *25th USENIX Security Symposium (USENIX Security 16)*, 601–18.
- Tschand, Arya, Arun Tejusve Raghunath Rajan, Sachin Idgunji, Anirban Ghosh, Jeremy Holleman, Csaba Kiraly, Pawan Ambalkar, et al. 2024. "MLPerf Power: Benchmarking the Energy Efficiency of Machine Learning Systems from μ Watts to MWatts for Sustainable AI." *2025 IEEE International Symposium on High Performance Computer Architecture (HPCA)*, 1201–16. <https://doi.org/10.1109/hpca61900.2025.00092>.
- Turakhia, Yatish, Gill Bejerano, and William J. Dally. 2018. "Darwin: A Genomics Co-Processor Provides up to 15,000X Acceleration on Long Read Assembly." *Proceedings of the Twenty-Third International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, 199–213. <https://doi.org/10.1145/3173162.3173193>.
- Turing, Alan M. 1950. "I.—Computing Machinery and Intelligence." *Mind* LIX (236): 433–60. <https://doi.org/10.1093/mind/lix.236.433>.
- Ultralytics. 2023. *YOLOv8*.
- Umuroglu, Yaman, Nicholas J. Fraser, Giulio Gambardella, Michaela Blott, Philip Leong, Magnus Jahre, and Kees Vissers. 2017. "Finn: A Framework for Fast, Scalable Binarized Neural Network Inference." *Proceedings of the 2017 ACM/SIGDA International Symposium on Field-Programmable Gate Arrays*, 65–74. <https://doi.org/10.1145/3020078.3021744>.
- United States Congress. 1996. *Health Insurance Portability and Accountability Act of 1996*. Public Law 104-191.
- U.S. Department of Health and Human Services. 2003. "Summary of the HIPAA Privacy Rule."
- U.S. Department of Health and Human Services. 2005. "Summary of the HIPAA Security Rule."
- U.S. Department of Health and Human Services. 2026. *Annual Civil Monetary Penalties Inflation Adjustment*. Federal Register.
- U.S. Food and Drug Administration. 2021. *Artificial Intelligence/Machine Learning (AI/ML)-Based Software as a Medical Device (SaMD) Action Plan*. U.S. Department of Health; Human Services.
- U.S. Food and Drug Administration. 2025. *Marketing Submission Recommendations for a Predetermined Change Control Plan for Artificial Intelligence-Enabled Device Software Functions*. Guidance for Industry and Food and Drug Administration Staff. U.S. Department of Health; Human Services.
- U.S. Securities and Exchange Commission. 2013. "Securities Exchange Act of 1934, Release No. 70694: Knight Capital Americas LLC."
- Vasisht, Deepak, Zerina Kapetanovic, Jongho Won, Xinxin Jin, Ranveer Chandra, Sudipta N. Sinha, Ashish Kapoor, Madhusudhan Sudarshan, and Sean Stratman. 2017. "FarmBeats: An IoT Platform for Data-Driven Agriculture." *14th USENIX Symposium on Networked Systems Design and Implementation (NSDI 17)*, 515–29.
- Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. "Attention Is All You Need." *Advances in Neural Information Processing Systems (NeurIPS)* 30: 5998–6008.
- Villalobos, Pablo, Anson Ho, Jaime Sevilla, Tamay Besiroglu, Lennart Heim, and Marius Hobbhahn. 2022. "Will We Run Out of Data? Limits of LLM Scaling Based on Human-Generated Data." *arXiv Preprint arXiv:2211.04325*.
- Viola, Paul, and Michael J. Jones. 2001. "Rapid Object Detection Using a Boosted Cascade of Simple Features." *Proceedings of the 2001 IEEE Computer Society Conference on Computer Vision and Pattern Recognition. CVPR 2001* 1: 1-511-1-518. <https://doi.org/10.1109/cvpr.2001.990517>.
- Wang, Alex, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R. Bowman. 2019. "SuperGLUE: A Stickier Benchmark for General-Purpose Language Understanding Systems." *arXiv Preprint arXiv:1905.00537*.

- Wang, Alex, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R. Bowman. 2018. "GLUE: A Multi-Task Benchmark and Analysis Platform for Natural Language Understanding." *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, 353–55. <https://doi.org/10.18653/v1/w18-5446>.
- Wang, Linnan, Jinmian Ye, Yiyang Zhao, Wei Wu, Ang Li, Shuaiwen Leon Song, Zenglin Xu, and Tim Kraska. 2018. "SuperNeurons: Dynamic GPU Memory Management for Training Deep Neural Networks." *Proceedings of the 23rd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, 41–53. <https://doi.org/10.1145/3178487.3178491>.
- Wang, Shibo, and Pankaj Kanwar. 2019. *BFloat16: The Secret to High Performance on Cloud TPUs*.
- Wang, Tianlu, Jieyu Zhao, Mark Yatskar, Kai-Wei Chang, and Vicente Ordonez. 2019. "Balanced Datasets Are Not Enough: Estimating and Mitigating Gender Bias in Deep Image Representations." *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, 5309–18. <https://doi.org/10.1109/iccv.2019.00541>.
- Wang, Xin, Fisher Yu, Zi-Yi Dou, Trevor Darrell, and Joseph E. Gonzalez. 2018. "SkipNet: Learning Dynamic Routing in Convolutional Networks." In *European Conference on Computer Vision (ECCV)*. Springer. https://doi.org/10.1007/978-3-030-01261-8_25.
- Wang, Yu Emma, Gu-Yeon Wei, and David Brooks. 2019. "Benchmarking TPU, GPU, and CPU Platforms for Deep Learning." *arXiv Preprint arXiv:1907.10701*.
- Wang, Zijie J., Robert Turko, Omar Shaikh, Haekyu Park, Nilaksh Das, Fred Hohman, Minsuk Kahng, and Duen Horng Polo Chau. 2021. "CNN Explainer: Learning Convolutional Neural Networks with Interactive Visualization." *IEEE Transactions on Visualization and Computer Graphics* 27 (2): 1396–406. <https://doi.org/10.1109/tvcg.2020.3030418>.
- Warden, Pete. 2018. "Speech Commands: A Dataset for Limited-Vocabulary Speech Recognition." *arXiv Preprint arXiv:1804.03209*, ahead of print. <https://doi.org/10.48550/arXiv.1804.03209>.
- Warden, Pete, and Daniel Situnayake. 2020. *TinyML: Machine Learning with TensorFlow Lite on Arduino and Ultra-Low-Power Microcontrollers*. O'Reilly Media.
- Waterman, Andrew, Yunsup Lee, Rimas Avizienis, Henry Cook, David Patterson, and Krste Asanovic. 2013. "The RISC-V Instruction Set." *2013 IEEE Hot Chips 25 Symposium (HCS)*, 1–1. <https://doi.org/10.1109/hotchips.2013.7478332>.
- Wei, Jason, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. 2022. "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models." *Advances in Neural Information Processing Systems (NeurIPS)* 35: 24824–37. <https://doi.org/10.52202/068431-1800>.
- Weights & Biases, Inc. 2024. *Weights & Biases: The AI Developer Platform*.
- Weiser, Mark. 1991. "The Computer for the 21st Century." *Scientific American* 265 (3): 94–104. <https://doi.org/10.1038/scientificamerican0991-94>.
- Weizenbaum, Joseph. 1966. "ELIZA—a Computer Program for the Study of Natural Language Communication Between Man and Machine." *Communications of the ACM* 9 (1): 36–45. <https://doi.org/10.1145/365153.365168>.
- Welling, Max. 2009. "Herding Dynamical Weights to Learn." *Proceedings of the 26th Annual International Conference on Machine Learning*, 1121–28. <https://doi.org/10.1145/1553374.1553517>.
- Wengert, Ralph E. 1964. "A Simple Automatic Derivative Evaluation Program." *Communications of the ACM* 7 (8): 463–64. <https://doi.org/10.1145/355586.364791>.
- Werbos, Paul. 1974. "Beyond Regression: New Tools for Prediction and Analysis in the Behavioral Sciences." PhD thesis, Harvard University.
- Werchaniak, Andrew, Roberto Barra Chicote, Yuriy Mishchenko, Jasha Droppo, Jeff Condal, Peng Liu, and Anish Shah. 2021. "Exploring the Application of Synthetic Audio in Training Keyword Spotters." *ICASSP 2021 - 2021 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 7993–96. <https://doi.org/10.1109/icassp39728.2021.9413448>.
- Wexler, James, Mahima Pushkarna, Tolga Bolukbasi, Martin Wattenberg, Fernanda Viégas, and Jimbo Wilson. 2020. "The What-If Tool: Interactive Probing of Machine Learning Models." *IEEE Transactions on Visualization and Computer Graphics* 26 (1): 56–65. <https://doi.org/10.1109/TVCG.2019.2934619>.
- Widmer, Gerhard, and Miroslav Kubat. 1996. "Learning in the Presence of Concept Drift and Hidden Contexts." *Machine Learning* 23 (1): 69–101. <https://doi.org/10.1023/a:1018046501280>.
- Williams, Samuel, Andrew Waterman, and David Patterson. 2009. "Roofline: An Insightful Visual Performance Model for Multicore Architectures." *Communications of the ACM* 52 (4): 65–76. <https://doi.org/10.1145/1498765.1498785>.
- Wolpert, D. H., and W. G. Macready. 1997. "No Free Lunch Theorems for Optimization." *IEEE Transactions on Evolutionary Computation* 1 (1): 67–82. <https://doi.org/10.1109/4235.585893>.
- Wu, Bichen, Kurt Keutzer, Xiaoliang Dai, Peizhao Zhang, Yanghan Wang, Fei Sun, Yiming Wu, Yuandong Tian, Peter Vajda, and Yanguang Jia. 2019. "FBNet: Hardware-Aware Efficient ConvNet Design via Differentiable Neural Architecture Search." *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 10726–34. <https://doi.org/10.1109/cvpr.2019.01099>.

- Wu, Hao, Patrick Judd, Xiaojie Zhang, Mikhail Isaev, and Paulius Micikevicius. 2020. "Integer Quantization for Deep Learning Inference: Principles and Empirical Evaluation." *arXiv Preprint arXiv:2004.09602* abs/2004.09602.
- Wu, Jian, Hao Cheng, and Yifan Zhang. 2019. "Fast Neural Networks: Efficient and Adaptive Computation for Inference." *Advances in Neural Information Processing Systems (NeurIPS)*.
- Wulf, Wm. A., and Sally A. McKee. 1995. "Hitting the Memory Wall: Implications of the Obvious." *ACM SIGARCH Computer Architecture News* 23 (1): 20–24. <https://doi.org/10.1145/216585.216588>.
- Xin, Ji, Raphael Tang, Yaoliang Yu, and Jimmy Lin. 2021. "BERxIT: Early Exiting for BERT with Better Fine-Tuning and Extension to Regression." In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, edited by Paola Merlo, Jorg Tiedemann, and Reut Tsarfaty. Association for Computational Linguistics. <https://doi.org/10.18653/v1/2021.eacl-main.8>.
- Xu, Ruijie, Zengzhi Wang, Run-Ze Fan, and Pengfei Liu. 2024. "Benchmarking Benchmark Leakage in Large Language Models." *arXiv Preprint arXiv:2404.18824*.
- Yang, Le, Yizeng Han, Xi Chen, Shiji Song, Jifeng Dai, and Gao Huang. 2020. "Resolution Adaptive Networks for Efficient Inference." *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2366–75. <https://doi.org/10.1109/cvpr42600.2020.00244>.
- Yao, Zhewei, Amir Gholami, Sheng Shen, Kurt Keutzer, and Michael W. Mahoney. 2021. "HAWQ-V3: Dyadic Neural Network Quantization." *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 11875–86.
- Yee, Kyra, Uthaiapon Tantipongpipat, and Shubhanshu Mishra. 2021. "Image Cropping on Twitter: Fairness Metrics, Their Limitations, and the Importance of Representation, Design, and Agency." *Proceedings of the ACM on Human-Computer Interaction* 5 (CSCW2): 1–24. <https://doi.org/10.1145/3479594>.
- Yosinski, Jason, Jeff Clune, Yoshua Bengio, and Hod Lipson. 2014. "How Transferable Are Features in Deep Neural Networks?" *Advances in Neural Information Processing Systems* 27.
- Yu, Gyeong-In, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. 2022. "Orca: A Distributed Serving System for Transformer-Based Generative Models." *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*, 521–38.
- Yun, Sangdoo, Dongyoon Han, Seong Joon Oh, Sanghyuk Chun, Junsuk Choe, and Youngjoon Yoo. 2019. "CutMix: Regularization Strategy to Train Strong Classifiers with Localizable Features." *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, 6023–32. <https://doi.org/10.1109/ICCV.2019.00612>.
- Yurdakul, Bilal, and Joshua Naranjo. 2020. "Statistical Properties of the Population Stability Index." *The Journal of Risk Model Validation* 52. <https://doi.org/10.21314/jrmv.2020.227>.
- Zafir, Ofir, Guy Boudoukh, Peter Izsak, and Moshe Wasserblat. 2019. "Q8BERT: Quantized 8Bit BERT." *2019 Fifth Workshop on Energy Efficient Machine Learning and Cognitive Computing - NeurIPS Edition (EMC2-NIPS)*, 36–39. <https://doi.org/10.1109/emc2-nips53020.2019.00016>.
- Zaharia, Matei, Andrew Chen, Aaron Davidson, Ali Ghodsi, Sue Ann Hong, Andy Konwinski, Siddharth Murching, et al. 2018. "Accelerating the Machine Learning Lifecycle with MLflow." *IEEE Data Engineering Bulletin* 41 (4): 39–45.
- Zaharia, Matei, Mosharaf Chowdhury, Michael J. Franklin, Scott Shenker, and Ion Stoica. 2010. "Spark: Cluster Computing with Working Sets." *HotCloud* 10: 10–10.
- Zaharia, Matei, Ali Ghodsi, Reynold Xin, and Michael Armbrust. 2021. "Lakehouse: A New Generation of Open Platforms That Unify Data Warehousing and Advanced Analytics." *Proceedings of CIDR* 8.
- Zech, John R., Marcus A. Badgeley, Manway Liu, Anthony B. Costa, Joseph J. Titano, and Eric Karl Oermann. 2018. "Variable Generalization Performance of a Deep Learning Model to Detect Pneumonia in Chest Radiographs: A Cross-Sectional Study." *PLOS Medicine* 15 (11): e1002683. <https://doi.org/10.1371/journal.pmed.1002683>.
- Zhang, Biao, and Rico Sennrich. 2019. "Root Mean Square Layer Normalization." *Advances in Neural Information Processing Systems (NeurIPS)*.
- Zhang, Brian Hu, Blake Lemoine, and Margaret Mitchell. 2018. "Mitigating Unwanted Biases with Adversarial Learning." *Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society*, 335–40. <https://doi.org/10.1145/3278721.3278779>.
- Zhang, Chengliang, Minchen Yu, Wei Wang, and Feng Yan. 2019. "Mark: Exploiting Cloud Services for Cost-Effective, SLO-Aware Machine Learning Inference Serving." *2019 USENIX Annual Technical Conference (USENIX ATC 19)*, 1049–62.
- Zhang, Chiyuan, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. 2017. "Understanding Deep Learning Requires Rethinking Generalization." *International Conference on Learning Representations (ICLR)*.
- Zhang, Hongyi, Moustapha Cisse, Yann N. Dauphin, and David Lopez-Paz. 2018. *mixup: Beyond Empirical Risk Minimization*. OpenReview.net.

- Zhang, Li Lina, Yuqing Yang, Yuhang Jiang, Wenwu Zhu, and Yunxin Liu. 2020. "Fast Hardware-Aware Neural Architecture Search." *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*. <https://doi.org/10.1109/cvprw50498.2020.00354>.
- Zhang, Qingxue, Dian Zhou, and Xuan Zeng. 2017. "Highly Wearable Cuff-Less Blood Pressure and Heart Rate Monitoring with Single-Arm Electrocardiogram and Photoplethysmogram Signals." *BioMedical Engineering OnLine* 16 (1): 23. <https://doi.org/10.1186/s12938-017-0317-z>.
- Zhang, Yundong, Naveen Suda, Liangzhen Lai, and Vikas Chandra. 2017. *Hello Edge: Keyword Spotting on Microcontrollers*.
- Zhao, Jiawei, Zhenyu Zhang, Beidi Chen, Zhangyang Wang, Anima Anandkumar, and Yuandong Tian. 2024. "GaLore: Memory-Efficient LLM Training by Gradient Low-Rank Projection." *arXiv Preprint*.
- Zheng, Lianmin, Chengfan Jia, Minmin Sun, Zhao Wu, Cody Hao Yu, Ameer Haj-Ali, Yida Wang, et al. 2020. "Ansor: Generating High-Performance Tensor Programs for Deep Learning." *Proceedings of the 14th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 863–79.
- Zhou, Aojun, Yukun Ma, Junnan Zhu, Jianbo Liu, Zhijie Zhang, Kun Yuan, Wenxiu Sun, and Hongsheng Li. 2021. "Learning n:m Fine-Grained Structured Sparse Neural Networks from Scratch." *arXiv Preprint*.
- Zhu, Chenzhao, Song Han, Huizi Mao, and William J. Dally. 2017. "Trained Ternary Quantization." *International Conference on Learning Representations (ICLR)* 4 (4): 1–16.
- Zhu, Ligeng, Zhijian Liu, and Song Han. 2019. "Deep Leakage from Gradients." *Advances in Neural Information Processing Systems (NeurIPS)* 32: 14774–84.
- Zillow Group. 2021. *Zillow Group Reports Third-Quarter 2021 Financial Results and Shares Plan to Wind down Zillow Offers Operations*. Investor Relations Press Release.
- Zoph, Barret, and Quoc V. Le. 2016. "Neural Architecture Search with Reinforcement Learning." *International Conference on Learning Representations* 3.
- Zu, Y., A. Ghaffarkhah, H.-V. Dang, B. Towles, S. Hand, S. Huda, A. Bello, et al. 2024. "Resiliency at Scale: Managing Google's TPUs Machine Learning Supercomputer." *21st USENIX Symposium on Networked Systems Design and Implementation (NSDI 24)*, 761–74.

Index

- 1980s
 - definition, 538
- 1990s
 - definition, 538
- 2000s
 - definition, 538
- 2010s
 - definition, 538
- 2020s
 - definition, 538
- A/B testing
 - delayed feedback, A/B testing challenges, 766
 - ML deployment, 95
 - randomization unit, 765
 - statistical model validation, 765
- A100, 379, 391
 - data-loading requirements, 130
 - HBM capacity, 344
 - MIG partitioning, 710
 - roofline analysis, 622, 891
 - Tensor Core throughput, 328
- Ablation study
 - component isolation, 94
 - definition, 94
- Abstraction problem
 - definition, 318
- Acceleration wall
 - diminishing returns, 530
- Accelerator
 - bubbles, pipeline idleness, 353
 - design, workload-specific memory, 571
 - memory bandwidth limitations, 370
 - ML, 537
 - specifications, roofline analysis example, 622
- Accelerator gain
 - definition, 529
- Access pattern
 - definition, 563
- Activation checkpointing
 - definition, 310
 - memory-compute trade-off, 310, 353, 399
 - optimal placement, 399
 - recomputation strategy, 312
 - recomputation trade-off, 379
- Activation function
 - definition, 185
 - nonlinear transformation, 178
 - softmax, 187
 - training efficiency, 360
- Activation memory
 - definition, 205
 - forward pass storage, 368
 - storage requirements, 368
- Activation reuse
 - definition, 571
- Activation unit
 - definition, 547
- Activation-based pruning
 - definition, 471
- Active learning
 - budget multiplier, 146
 - definition, 433
 - label prioritization, 146
 - labeling cost reduction, 425, 433
- Adam optimization, *see* Adam optimizer
- Adam optimizer
 - convergence properties, 364
 - Kingma and Ba, 364
 - memory overhead, 209, 364, 365
 - momentum and velocity, 209
- AdamW
 - decoupled weight decay, 365
- Adaptability
 - definition, 603
- Adaptive computation
 - early exit, 509
- Adaptive inference
 - definition, 509
- Adaptive resolution
 - content-based selection, 699
- Admission control
 - queue depth threshold, 706
- Adversarial debiasing
 - fairness constraints, 820
 - training cost, 820
- Ahead-of-time compilation, 301
 - predeployment, 729
- Ahead-of-time optimization, 296
- AI compute primitive
 - definition, 540
- AI democratization, 850
- AI engineering
 - definition, 24
- AI memory systems
 - bandwidth bottleneck, 559
- AI Memory Wall
 - definition, 559
- AI runtime
 - definition, 603
 - dynamic execution management, 603
 - vs. traditional runtime, 603
- AI winters
 - systems failure, 8
- AI-assisted labeling
 - preannotation, 146
 - scalability, 145
- Alert fatigue
 - operational risk, 784
- AlexNet, 264
 - benchmarking paradigm, 667
 - GPU deep learning era, 534
 - ImageNet breakthrough, 6, 175
 - ImageNet revolution, 237
 - memory constraint, 175
 - multi-GPU origin, 406
 - parameter growth, 152
 - two-GPU training, 534
- Algorithm-hardware adoption lag, 180
- Algorithmic efficiency
 - definition, 22
- Algorithmic fairness
 - healthcare ML, 89
- Alignment gap
 - proxy metric divergence, 808
- AllGather
 - collective communication, 327
- AllReduce
 - collective communication, 327
 - definition, 606
 - gradient synchronization, 409
 - gradient synchronization cost, 606
 - network wall, 409
 - ring topology, 409
- AlphaFold
 - cloud deployment, 27
- AlphaGo, 11
 - self-play training, 11
- Always-on sensing
 - power constraints, 45
- Amazon Kinesis
 - streaming trade-offs, 134
- Amazon recruiting
 - specification failure, 804
- Amdahl's law
 - acceleration ceiling, 530
 - appendix refresher, 892
 - data pipeline ceiling, 141
 - diagnostic analogy, 18
 - distributed training limit, 606
 - multi-accelerator scaling, 606
 - optimization ceiling, 650
 - parallel fraction, 529
 - pipeline bottleneck, 76
 - preprocessing bottleneck, 696, 846
 - speedup limits, 76
 - system optimization, 78
- Amdahl's reality
 - definition, 530
- Amdahl, Gene, *see also* Amdahl's law
- Anti-curriculum
 - hard examples first, 433
- AOT, *see* Ahead-of-time optimization
- Apache Airflow
 - lineage capture, 834
 - pipeline orchestration, 753
- Apache Atlas
 - data lineage, 834
- Apache Beam
 - distributed processing, 141
- Apache Kafka
 - event streaming, 134
 - ordering guarantee, 134
- Apache Pulsar
 - geo-replication, 134
- Apache Spark
 - distributed processing, 141
- Apache TVM
 - accelerator compilation, 578
 - ML compiler, 339
- Application scenario
 - definition, 25
- Application-specific integrated circuit, *see* ASIC
- Architectural efficiency
 - theory-practice gap, 504
- Architectural Integration
 - execution unit organization, 556
- Architecture
 - building blocks, 263
 - computational graph, 226
 - compute vs. memory trade-off, 230

- data-to-architecture mapping, 277
- encoder-decoder, 269
- parameter efficiency, 232
- representational power vs. efficiency, 227
- selection as bias matching, 232
- systems engineering, 226
- Architecture examples
 - definition, 538
- Area under the learning curve
 - definition, 455
- Area under the ROC curve
 - threshold independence, 93
- Argmax
 - prediction selection, 701
- Arithmetic intensity, 48
 - appendix refresher, 862, 874, 890
 - attention layer, 371
 - bottleneck diagnosis, 846
 - definition, 572
 - FLOP/byte ratio, 229
 - GPU utilization, 221
 - high intensity workloads, 45
 - roofline diagnostic, 531
 - threshold for compute vs. memory bound, 622
 - training operations, 370
- Artifact
 - interdependence, 95
 - lineage tracking, 95
 - reproducibility, 742
 - versioning, reproducibility requirement, 742
- Artificial intelligence
 - hierarchy with ML and DL, 171
- Artificial neuron
 - computing primitive, 178
- ASIC
 - benchmarking trade-off, 622
 - flexibility trade-off, 536
- Assumptions
 - accelerator, 900
 - bandwidth, 906
 - energy, 905
 - nameplate capacity, 908
 - pricing, 907
 - reference models, 904
 - scale, 908
 - training memory, 904
 - unit conventions, 908
- Asynchronous execution
 - hiding transfer latency, 322
- Attention
 - memory bottleneck, 394
 - quadratic complexity, 394
- Attention head
 - specialization, 257
- Attention matrix, 256
- Attention Mechanism
 - global token interaction, 570
 - memory-bound softmax, 574
- Attention mechanism
 - Bahdanau, 250
 - content-addressable memory, 257
 - definition, 250
 - fused kernel, 292
 - information bottleneck, 250
 - memory pressure, 570
 - multi-head, 257
 - quadratic bottleneck, 255
 - query-key-value, 252
 - scaled dot-product, 257
 - self-attention, 257
 - similarity computation, 378
- AUC
 - anomaly detection performance, 633
- Audit trail
 - accountability logging, 834
 - inference-time logging, 835
 - multi-tier logging, 834
 - storage scale, 836
- AutoAugment
 - search cost, 440
- autocast
 - mixed-precision context, 311
- Autograd
 - automatic differentiation, 209
 - engine, internals, 311
 - execution trade-off, 209
 - reverse-linked graph, 311
- Autograd tape
 - definition, 294
 - dynamic graph construction, 294
 - memory cost, 294
 - reverse traversal, 344
- Automatic differentiation
 - appendix refresher, 877
 - computational graph, 209, 367
 - definition, 306
 - forward mode, 307
 - forward vs. reverse mode, 307
 - forward-mode dual numbers, 307
 - framework role, 288
 - reverse mode, 205, 306, 308
 - reverse-mode backpropagation, 308
 - reverse-mode memory cost, 306
 - tape-based vs. transform-based, 316
- Automatic mixed precision
 - framework support, 389
- Automotive AI
 - real-time safety requirements, 608
- Autonomous perception
 - definition, 26
- Autonomous Vehicle
 - latency requirements, 60
- Autoregressive generation, 228
 - bandwidth bottleneck, 228
 - memory-bound, 45
- Autoregressive model
 - autoregressive, serial bottleneck, 722
 - token generation, 717
- Autoscaling
 - ML cold-start trade-off, 771
 - ML workloads, 771
 - startup latency, 731
- Back-translation
 - pipeline latency, 440
 - text augmentation, 439
- Backpressure
 - streaming systems, 132
- Backpropagation
 - activation storage, 264
 - appendix refresher, 877
 - computational asymmetry, 308
 - computational cost, 367
 - definition, 204
 - gradient computation, 204
 - historical development, 180
 - historical introduction, 357
 - memory cost, 180
 - memory requirements, 367, 536
 - memory trade-off, 264
 - provenance, 357
 - through time (BPTT), 248
- Backward pass
 - gradient computation, 379
 - memory operations, 379
 - triggering gradient computation, 344
- Bagging
 - ensemble method, 93
- Bandwidth
 - bottleneck, video streaming, 57
 - latency trade-off, 58, 895
 - PCIe, model swapping, 710
- Bandwidth hog
 - definition, 45
- Bandwidth taper
 - interconnect hierarchy, 557
- Baseline model
 - reference point selection, 631
- Batch
 - loss averaging, 202
 - training iteration, 197
- Batch inference
 - definition, 772
- Batch ingestion
 - definition, 131
- Batch normalization, 267
 - hardware implementation, 545
 - mode-dependent behavior, 331
 - systems cost, 187
 - training speedup, 267
 - training stability, 187
 - training-serving skew, 267
- Batch processing
 - definition, 363
 - hardware utilization, 197
- Batch Size
 - arithmetic intensity impact, 576
- Batch size, 236
 - arithmetic intensity effect, 372
 - gradient stability, 211
 - hardware alignment, 377
 - inference trade-off, 215
 - loss computation, 202
- Batch throughput
 - definition, 650
- Batch-1 regime
 - definition, 716
- Batch-N regime
 - definition, 716
- Batching
 - batch formation delay, definition, 703
 - batch, small, CPU advantage, 729
 - etymology, 711
 - fallacy of larger batches, 735
 - hardware utilization, 359
 - request, throughput optimization, 769
 - tax, latency penalty, 703
 - training vs. serving, 711
 - window, latency-throughput trade-off, 694
- Battery
 - coin-cell deployment, 65
- Battery life
 - ML impact, 62
- Battery tax
 - always-on constraints, 61, 62
- Bayes error rate
 - appendix refresher, 866
- Beam search
 - decoding strategy, 723
- Benchmark
 - etymology, 616
- Benchmark contamination
 - LLM memorization risk, 670
- Benchmark coverage
 - incomplete problem representation, 662
- Benchmark engineering
 - intentional optimization bias, 664
- Benchmark gaming
 - distorting performance metrics, 666
 - vendor optimization for tests, 616
- Benchmark harness
 - test infrastructure, 633
- Benchmark saturation
 - outdated benchmarks pitfall, 676
- Benchmark-production gap, 614
- Benchmarking
 - battery, duty cycle specification, 638
 - benchmarks as proxies, 615
 - definition, 614
 - end-to-end vs. component metrics, 616
 - energy, first-class metric, 617
 - fallacy of direct performance translation, 676
 - historical evolution, 616
 - probabilistic variability, 618
 - system, definition, 614
- Benchmarking granularity

- definition, 625
- Bengio, Yoshua
 - Turing Award, 171
- Berkson, Joseph, *see also* Logits
- BERT
 - benchmarking baseline, 631
 - compression techniques, 476
 - inference deployment
 - prediction, 623
 - roofline analysis at batch 1, 623
- Best when
 - definition, 594
- BF16
 - appendix refresher, 896
 - design rationale, 313
 - dynamic range preservation, 392
 - exponent range preservation, 388
 - hardware support, 392
 - reduced-precision format, 215
- Bias
 - activation threshold, 190
 - attribute removal fallacy, 836
 - gender discrimination, 805
 - geographic, training data, 122
 - historical data encoding, 805
 - mitigation strategies, 809
 - parameter overhead, 190
 - proxy variables, 805
 - signal aggregation, 178
- Bias feedback model
 - disparity amplification, 846
- Bias-variance tradeoff
 - classical theory, 175
 - double descent, 175
 - overparameterization, 175
- Binarization
 - definition, 503
- Biological Neuron
 - definition, 178
- Bisection bandwidth
 - fleet-scale bottleneck, 31
- Bitter lesson
 - definition, 10
- bitter lesson, The, 10, 11
- BLAS
 - historical foundation, 290
 - matrix computation foundation, 358
 - neural networks, 235
 - performance foundation, 290
 - training compute hierarchy, 358
 - utilization, 235
 - vector and matrix operations, 329
- BLEU
 - Goodhart's law example, 632
- Blue-green deployment
 - duplicate capacity, 765
 - zero-downtime updates, 764
- Bobrow, Daniel
 - STUDENT, 8
- Boosting
 - ensemble method, 93
- Bottleneck principle
 - definition, 43
- Boundary erosion
 - dissolution of modularity, 747
- Bounding box
 - annotation, 143
 - definition, 143
- Break-even analysis
 - edge vs. cloud utilization, 54
- Break-even point
 - data selection ROI, 449
- Brittleness
 - definition, 10
- Broadcasting
 - definition, 876
- Bulkhead pattern
 - failure containment, 791
 - resource isolation, 791
- Byte pair encoding
 - tokenization, 375
- Byzantine fault tolerance
 - ML-specific manifestations, 791
 - plausible ML failures, 791
- C4
 - quality filtering, 428
- CACE principle
 - change propagation, 748
- Cache efficiency
 - definition, 588
- Cache locality
 - definition, 563
- Calibration
 - clinical trust, 97
 - dataset, representative traffic, 728
 - definition, 97
 - definition and etymology, 667
 - INT8 quantization, 728
 - production traffic mismatch, 735
- Canary
 - etymology, 706
- Canary deployment
 - risk-controlled rollout, 764
 - statistical validation, 764
- Canary request
 - fan-out protection, 706
- Canonical workload, 229
- CAP theorem
 - feature store trade-off, 134
 - streaming systems, 134
- Capacity planning
 - infrastructure sizing, 732
- Carbon footprint
 - energy to emissions conversion, 828
- Carbon intensity
 - definition, 830
 - grid emissions, 830
- Carbon-aware scheduling
 - lower-carbon times and regions, 830
- Categorical Encoding
 - one-hot encoding, 137
- Categorical feature, 261
- Central processing unit, *see* CPU
- Chain Rule
 - gradient decomposition, 205
- Chain rule
 - appendix refresher, 877
 - backpropagation foundation, 367
 - depth constraint, 206
- Challenge
 - definition, 579
- Channelwise quantization, *see* Quantization
- Characteristics
 - definition, 538
- CheckList, 811
- Checkout-free retail
 - edge processing, 60
- Checkpointing
 - I/O cost, 213
- Chinchilla
 - compute-optimal scaling, 455
 - compute-optimal training, 455
 - tokens-per-parameter diagnostic, 456
- Chiplet
 - modular die interconnect, 606
- CI/CD pipeline
 - ML-specific adaptation, 756
- CIFAR-10
 - classification reference dataset, 631
- Circuit breaker
 - cascade prevention, 130
- Circuit breaker pattern
 - cascade failure prevention, 780
 - ML adaptation, 780
- Class balance
 - coverage metric, 671
- Class imbalance
 - coverage failure mode, 671
 - minority-class performance, 221
- Classical machine learning
 - feature engineering, 173
- Classification
 - definition, 573
- Classification label
 - storage requirements, 143
- Cleanlab
 - label error detection, 432
- ClinAIops
 - healthcare ML operations, 795
 - MLOps comparison, 797
- Clinician oversight loop, 795
- CLIP, *see* Contrastive Language-Image Pretraining
- Closed division
 - MLPerf submission, 654
- Cloud economics
 - elastic pricing, 53
- Cloud infrastructure
 - utility model, 50
- Cloud ML, 26
 - accelerator infrastructure, 52
 - characteristics, 40
 - consumer-scale applications, 54
 - data center scale, 50
 - definition, 51
 - latency limitations, 53
 - privacy concerns, 53
 - scaling limits, 54
 - workload archetypes, 50
- Cloud pricing
 - cost-latency trade-off, 120
- CNN
 - definition, 236
 - spatial data reuse, 570
- COCO
 - object detection dataset, 631
- Code generation
 - hardware-specific instructions, 602
- Codec
 - definition, 536
- Cohen's kappa
 - inter-annotator agreement, 671
 - label quality, 128
- Coin-cell battery
 - deployment longevity, 64
- Cold start
 - anatomy, 708
 - bursty traffic impact, 735
 - cloud serving, 690
 - definition, 708
 - inference concern, 653
 - scaling events, 708
- Cold start latency
 - physical lower bound, 653
 - serverless AI impact, 651
- Collaborative Filtering
 - recommendation systems, 55
- Collation
 - batch assembly, 326
- Columnar storage
 - Parquet, 154
- Common techniques
 - definition, 576
- Common use cases
 - definition, 594
- Communication and synchronization
 - definition, 580
- Communication tax
 - definition, 405
- COMPAS
 - calibration and error rates, 806
 - recidivism algorithm audit, 805
- Compilation continuum
 - principle, 302
- Compilation output
 - definition, 597
- Complexity tax
 - ML vs. heuristics, 71
- Compositional representation, 242
- Compositionality
 - pattern decomposition, 183
- Compressed sparse row
 - definition, 875
- Compression
 - algorithm selection, 155
- Computation graph
 - dynamic (eager execution), 209
 - static (deferred execution), 209
- Computation Placement

- processing element assignment, 578
- Computation scheduling
 - definition, 597
 - parallel execution, 601
- Computational granularity
 - definition, 580
- Computational graph
 - definition, 288
 - definition (DAG), 292
 - framework optimization, 288
 - static, 295
- Computational pattern
 - definition, 538
- Computational photography
 - pipeline constraint, 61
- Compute beast
 - definition, 45
- Compute bottleneck
 - definition, 10
- Compute efficiency
 - definition, 22
- Compute-Bound
 - roofline classification, 571
- Compute-bound
 - operation classification, 291
 - resolution scaling, 699
- Compute-Bound vs. Memory-Bound
 - training vs. inference, 48
- Compute-bound vs. memory-bound
 - definition, 43
 - memory-bound optimization, 44
- Compute-optimal frontier
 - Chinchilla scaling laws, 455
- Computer engineering
 - historical lineage, 24
- Concept drift
 - adversarial evolution, 745
 - changing relationships, 773
 - definition, 128, 806
 - validation, 98
- Concept-based explanation
 - human-understandable feature, 822
- Conditional computation
 - definition, 509
- Confidence interval
 - benchmark reporting, 619
- Confidence threshold
 - automation trade-off, 217
 - decision making, 218
- Confident learning
 - mislabeled example detection, 671
- Configuration debt
 - Facebook case study, 751
 - parameter sprawl, 750
- Confusion matrix
 - fairness analysis, 818
- Consensus labeling
 - quality control, 145
- Conservation of complexity
 - heuristic limits, 845
 - meta-principle, 845
 - structural optimization, 467
- Consistency Imperative
 - definition, 136
- Consistency imperative
 - training-serving parity, 136, 743
- Consistency regularization
 - compute trade-off, 436
 - semi-supervised technique, 436
- Constant folding, 297
 - compile-time optimization, 728
- Constrained hardware
 - inference deployment, 214
- Constraint layer
 - statistical physical operational, 89
- Constraint propagation principle
 - definition, 102
- Containerization
 - environment parity, 764
 - ML deployment, 770
 - reproducible environments, 142
- Continuous batching
 - LLM serving, 717
 - scheduling granularity, 717
- Continuous therapeutic monitoring
 - healthcare AI, 794
 - human oversight, 794
- Contrastive Language-Image Pretraining
 - semantic deduplication, 431
- Contrastive Learning
 - batch-based methods, 439
 - pretext task, 437
- Convergence
 - monitoring, 212
 - rate, clean vs. noisy data, 428
 - time-memory trade-off, 365
- Convolution
 - computational cost, 378
 - data reuse, 236
 - depthwise, 245
 - depthwise separable, 245
 - etymology, 236
 - hardware optimization, 245
 - kernel, 238
 - matrix multiplication
 - equivalence, 543
 - padding, 241
 - pointwise, 245
 - stride, 241
- Convolution performance
 - definition, 588
- Convolutional neural network, *see* CNN
- Coordination Tax
 - distributed data selection, 452
- Coordination tax
 - definition, 140
- Coreset
 - data efficiency, 443
 - definition, 429
 - etymology, 429
 - geometry origins, 429
- Correction cascade
 - sequential dependencies, 748
 - Zillow case study, 751
- Cosine annealing
 - learning rate schedule, 367
- Cost amortization
 - self-supervised pretraining, 438
- Cost modeling
 - data selection economics, 447
- Cost per inference
 - calculation, 772
 - serving economics, 731
- Cost-aware automation
 - retraining decisions, 743
- Cost-Performance Analysis
 - accelerator economics, 557
- Counterfactual analysis
 - prediction debugging, 783
- Covariate shift
 - definition, 128
 - input distribution change, 672
- CP decomposition
 - definition, 480
- CPU
 - inference optimization, 729
 - neural network execution, 201
- CPU affinity
 - latency reduction, 693
- CPU backend
 - optimized BLAS libraries, 346
- Cray-1
 - AI accelerator lineage, 542
- Credit assignment problem
 - weight responsibility, 204
- Critical batch size
 - throughput-convergence limit, 363
- Critical path
 - composed request, 852
 - optimization, 47
- Cross-entropy loss
 - classification loss, 202
 - definition, 202
 - information theory, 202
- Cross-industry standard process for data mining
 - lifecycle influence, 80
- Cross-layer interaction
 - cross-optimization effects, 452
- Crowdsourcing
 - annotation, 121
 - data labeling, 121
- cuBLAS
 - hardware-accelerated linear algebra, 358
 - kernel selection, 328
 - NVIDIA linear algebra, 599
- CUDA
 - context, initialization overhead, 708
 - cores, streaming processors, 578
 - dispatch overhead, 300
 - ecosystem lock-in, 548
 - runtime, OS layer, 321
 - serving cold starts, 708
 - stream definition, 322
 - stream overlap, 322
- CUDA graphs
 - kernel launch replay, 627
- CUDA MPS
 - context sharing, 708
 - multi-model serving, 708
- cuDNN
 - benchmarking dependency, 626
 - NVIDIA deep learning library, 599
 - optimized implementations, 378
- Curriculum learning
 - difficulty scorer and pacing function, 432
 - etymology, 432
 - optimization landscape, 432
- Custom autograd function
 - backward implementation, 312
- Custom differentiation rule, 312
- CutMix
 - image augmentation, 439
 - label mixing, 440
- Cutout
 - augmentation efficiency, 439
 - image augmentation, 439
- D-A-M Taxonomy
 - bottleneck classification, 382
- DAM taxonomy, *see* D-A-M taxonomy
- D-A-M taxonomy
 - alignment of components, 220
 - appendix refresher, 859
 - axis intersections, 860
 - definition, 13
 - diagnostic framework, 674
 - responsibility diagnosis, 808
 - scorecard, 865
 - training vs. inference, 48
 - workload classification, 45
- Dark knowledge
 - class correlations, 478
- Dark Silicon
 - specialization driver, 535
- Dartmouth Conference, 8
- systems oversight, 8
- Data
 - as source code, 13
 - coverage, requirements, 122
 - partitioning, query optimization, 155
 - physics of, 109
 - transformation, techniques, 137
- Data acquisition
 - scalability, 119
 - scale vs. quality, 119
 - source evaluation, 118
 - strategies, 118
- Data augmentation
 - definition, 439
 - synthetic data, 121
 - training diversity, 217
- Data benchmarking
 - definition, 615
 - training set validation, 670
- Data card
 - compliance documentation, 834
- Data cascade

- definition, 102, 112
- silent failure, 111
- Data center
 - energy consumption, 830
 - ML infrastructure, 50
- Data cleaning
 - error correction, 136
 - inconsistency removal, 136
 - signal-to-noise engineering, 110
- Data collection
 - iron law data term, 90
- Data compression ratio
 - definition, 455
- Data curation
 - end-to-end generalization, 625
- Data debt
 - definition, 159
 - hidden liability, 158
 - strategic vs. unconscious, 160
- Data Drift
 - definition, 773, 806
- Data drift
 - definition, 17
 - degradation equation, 137
 - detection, 127
 - operational burden, 28
 - silent model degradation, 100
 - Waymo example, 28
- Data echoing
 - amortizing I/O costs, 446
 - GPU utilization, 446
- Data engineering
 - definition, 29, 108
- Data freshness
 - SLA, 132
 - staleness detection, 778
- Data governance
 - definition, 831
 - enforcement mechanisms, 831
 - model artifact compliance, 837
- Data gravity
 - appendix refresher, 867
 - definition, 109
- Data gravity invariant
 - cloud limitations, 53
- Data gravity principle
 - physical anchor, 845
- Data ingestion
 - pipeline boundary, 130
- Data labeling
 - clinical economics, 436
 - systems engineering, 143
- Data lake
 - schema flexibility, 150
 - unstructured storage, 150
- Data lakehouse
 - architecture, 109
 - data gravity, 109
- Data lineage
 - audit trail, 102
 - debugging and compliance, 142
 - GDPR compliance, 834
 - processing history, 142
 - regulatory compliance, 101
 - transformation tracking, 142
- Data loading
 - shard-based sequential I/O, 446
- Data Locality
 - spatial and temporal, 552
- Data locality
 - code-to-data, 111
 - data-to-code, 111
- Data locality invariant
 - definition, 58
- Data localization
 - cross-border transfer, 834
- Data management
 - MLOps lifecycle, 752
- Data mesh
 - decentralized ownership, 109
- Data minimization
 - privacy by design, 834
- Data movement
 - broadcast, 274
 - energy dominance, 42, 177, 583
 - gather, 274
 - optimization strategies, 583
 - reduction, 274
 - scatter, 274
- Data Parallelism
 - training strategy, 642
- Data parallelism
 - communication overhead, 356
 - framework implementation, 406
 - gradient averaging, 406
 - gradient synchronization, 406
 - multi-accelerator scaling, 606
- Data path
 - customization, 535
- Data pipeline
 - architecture, 123
 - automated workflows, 753
 - ingestion and preprocessing, 373
 - storage throughput, 374
 - throughput optimization, 325
- Data poisoning
 - training pipeline security, 833
- Data prefetching
 - accelerator utilization, 385
 - buffer management, 385
- Data quality
 - as code, 126
 - container, 126
 - content, 126
 - expectation tests, 126
 - mechanical vs. semantic, 126
 - monitoring, 776
 - noise penalty on convergence, 428
 - statistical monitoring, 125
 - vs. data quantity, 77
- Data quality multiplier
 - label noise penalty, 429
- Data redundancy
 - empirical evidence, 428
- Data reuse
 - accelerator optimization, 584
 - definition, 563
 - energy efficiency, 177
- Data science
 - vs. data engineering, 108
- Data selection
 - definition, 22, 422, 424
 - optimization ordering, 423
 - systems vs. ML framing, 424
- Data shuffling
 - random access penalty, 374
- Data supply rate
 - bandwidth constraint, 154
- Data swamp
 - governance failure, 150
- Data tax
 - definition, 426
- Data validation
 - monitoring, 124
- Data versioning
 - DVC tool, 752
 - model reproducibility, 156
 - reproducibility, 83, 755
- Data wall
 - definition, 422
 - scaling constraint, 422
 - zero learning signal, 426
- Data warehouse
 - analytical queries, 149
 - OLAP, 149
- Data-as-code principle
 - definition, 5
 - logical program, 845
- Data-centric AI, 5
 - dataset enhancement paradigm, 672
- Data-model interface
 - feature consistency, 741
- Database
 - OLTP for ML, 149
- DataComp
 - data-centric benchmark, 672
 - data-centric benchmarking, 672
- Dataflow
 - optimization strategies, 584
- Dataflow Architecture
 - definition, 578
- Dataflow optimization
 - data movement minimization, 582
- Dataflow strategy
 - definition, 595
- DataHub
 - data lineage, 834
- DataLoader
 - configuration parameters, 387
 - throughput optimization, 325
- DataPerf
 - data quality benchmark, 672
- Dataset
 - iterable-style, 326
 - map-style, 326
 - MNIST benchmark, 191
- Dataset compilation
 - definition, 108
- Dataset saturation
 - performance capability illusion, 672
- Datasheets for datasets
 - training data documentation, 816
- DAWNBench
 - time-to-accuracy evaluation, 617
- Dead code elimination, 297
 - graph optimization, 297
- Dead neuron, 265
- Dean, Jeff
 - latency numbers, 151, 887
- Debugging
 - data pipelines, 161
- Debugging checklist
 - runbook steps, 783
- Decision framework
 - paradigm selection, 69
- Decode phase
 - memory-bandwidth bound, 734
 - sequential generation, 734
- Deduplication
 - data selection optimization, 431
 - definition, 108
 - embedding-based, 431
 - hash-based, exact matching, 431
 - storage and I/O cost reduction, 425
- Deep Blue, 11
 - custom chess hardware, 11
- Deep learning
 - automatic feature learning, 174
 - definition, 9, 171
 - network depth, 183
- Deep learning framework
 - automatic differentiation, 290
- Deep Learning Recommendation Model
 - workflow constraints, 85
- DeepBench
 - operator-level testing, 628
- Default in frameworks
 - definition, 588
- Deferred execution
 - graph construction, 288
- Degradation equation
 - definition, 17
 - distribution shift, 75
 - operational implementation, 740
- Delta Lake
 - time travel, 156
- Demographic drift
 - definition, 89
- Demographic parity
 - definition, 818
- Dennard Scaling
 - end of, 534
 - multi-core revolution, 536
- Dennard scaling
 - origin and breakdown, 42
 - power constraint, 535
- Dense connectivity, 233
 - parameter scaling, 192
- Dense Layer
 - roofline analysis, 575
- Deployment
 - competition-production gap, 93

- compute-intensive cloud
 - deployment, scientific computing, 27
 - controlled, risk mitigation, 762
 - high-stakes hybrid deployment, autonomous vehicles, 27
 - iron law overhead minimization, 98
 - resource-constrained edge, precision agriculture, 27
- Deployment infrastructure
 - definition, 29
- Deployment paradigm
 - definition, 15, 42
- Deployment target
 - cloud to edge spectrum, 339
- Depth
 - exponential advantage, 183
- Depthwise conv
 - definition, 573
- Depthwise Convolution
 - low arithmetic intensity, 573
- Depthwise separable convolution, 245
 - mobile deployment, 84
 - mobile efficiency, 505
 - MobileNet, 229
- Deterministic Transformation
 - reproducibility, 139
- Developer feedback loop, 795
- Device management
 - CPU-GPU transfers, 321
- DevOps
 - MLOps extension, 741
- Diabetic retinopathy
 - case study introduction, 86
 - deployment scale, 87
- Differential privacy
 - epsilon budget, 833
 - epsilon trade-off, 810
 - formal guarantees, 833
 - mathematical guarantees, 810
 - noise injection, 833
- Differentiation problem
 - definition, 306
- Diffusion model
 - text-to-image synthesis, 440
- Digital signal processing
 - multiply-accumulate units, 533
- Dimensional homogeneity
 - appendix refresher, 892
- Direct memory access, *see* DMA
- Disaggregated evaluation
 - definition, 811
 - demographic stratification, 811
 - fairness testing, 817
- Disparate impact
 - legal doctrine, 811
- Dispatch overhead
 - law, 304
- Dispatch overhead reduction
 - kernel launch batching, 627
- Dispatch tax
 - framework overhead, 295
- Display overlay
 - GPU mapping, 607
- Distance penalty
 - cloud latency, 51
- DistilBERT
 - deployment metrics, 478
- Distributed computing
 - system-level performance, 617
- Distributed data
 - clinic coordination cost, 92
- Distributed Processing
 - scaling, 140
- Distributed Selection
 - network topology constraints, 452
- Distributed training
 - communication overhead, 408
 - communication tax, 408
 - data parallelism, 356
 - data selection challenges, 450
 - execution contexts, 327
 - model parallelism, 356
 - scaling efficiency, 406
- Distribution Shift
 - responsibility lens, 806
- Distribution shift
 - causes of model degradation, 807
 - label shift, 128
 - monitoring, 127
 - population mismatch, 809
 - train-to-production alignment, 671
- Diversity Sampling
 - active learning query strategy, 434
- DLRM, 261
 - capacity-bound workload, 844
 - embedding bottleneck, 356
 - I/O-bound serving, 697
 - interaction layer, 261
 - memory capacity bound, 46
 - memory-bound benchmark, 654
 - memory-bound constraint, 55
 - recommendation model benchmark, 654
 - scalability challenge, 114
- DMA
 - computation overlap, 569
 - PCIe data movement interface, 321
- Documentation debt
 - data provenance, 159
- Domain adaptation
 - feature alignment, 441
- Domain gap
 - training-serving skew, 441
- Domain randomization
 - bridging synthetic-real gap, 440
 - rendering trade-off, 441
- Domain-Specific Architecture, *see* DSA
- Dot product
 - appendix refresher, 874
 - definition, 184
- Double buffering
 - latency hiding, 604
- Double descent
 - overparameterized regime, 175
- DQN
 - real-time inference, 172
- DRAM
 - energy cost, 538
- Drift
 - covariate shift, 773
- Drift detection delay
 - label latency, 773
 - statistical monitoring, 773
- Dropout
 - regularization technique, 187
 - train-inference mismatch, 187
 - training vs. inference, 331
- DS-CNN
 - monitoring strategy, 794
- DSA
 - definition, 534
 - efficiency trade-off, 534
 - scaling law breakdown, 534
- Dual mandate, 4
- Dual number
 - forward mode AD, 307
- Duty cycle
 - accelerator utilization, 700
 - ML lifecycle maintenance, 81
- DVC
 - data versioning, 745
 - dataset versioning, 752
 - git semantics, 156
- DVFS
 - dynamic voltage and frequency scaling, 659
 - power-performance envelope, 608
 - warm-up measurement artifact, 626
- Dying ReLU problem
 - definition, 187
- Dynabench
 - dynamic evaluation, 673
- Dynamic batching
 - definition, 711
- time window, 713
- Dynamic computation, 276
- Dynamic graph
 - definition, 294
 - eager execution, 293
 - runtime construction, 294
- Dynamic inference
 - real-time prediction, 689
- Dynamic kernel execution
 - runtime adaptation, 604
- Dynamic quantization
 - definition, 499
- Dynamic random access memory, *see* DRAM
- Dynamic selection
 - active learning, 433
 - curriculum learning, 432
 - definition, 423
 - training-time optimization, 432
- Dynamic voltage and frequency scaling, *see* DVFS
- Eager Execution
 - optimization limitations, 292
- Eager execution
 - define-by-run model, 293
 - definition, 288
 - optimization trade-off, 288
- Earth mover's distance
 - multi-dimensional drift, 774
- Edge accelerator
 - deployment selection, 59
- Edge AI
 - deployment patterns, 769
 - Oura Ring case study, 793
- Edge deployment
 - bandwidth constraints, 100
 - constraint triangle, 636
 - framework selection, 339
 - power constraints, 826
 - storage, 156
 - three problems reweighting, 339
 - TinyML frameworks, 305
- Edge hardware
 - resource limits, 58
- Edge inference
 - power tiers, 214
- Edge ML
 - autonomous vehicles, 60
 - bandwidth reduction, 57
 - data sovereignty, 56
 - definition, 56
 - deployment challenges, 57
 - distance penalty, 55
 - distributed processing, 57
 - industrial IoT, 60
 - latency and bandwidth, 26
 - latency benefits, 40
 - privacy benefits, 57
 - smart retail, 60
- Edge orchestration
 - fleet challenges, 58
- Edge serving
 - embedded function calls, 690
- Edge TPU
 - inference efficiency, 214
 - inference latency, 615
 - operator coverage, 646
- Effective Sample Size
 - label noise penalty, 429
- Efficiency
 - Jevons paradox, 24
- EfficientNet
 - algorithmic efficiency, 24
- Einstein summation convention
 - definition, 874
- EL2N
 - definition, 429
 - proxy transferability, 429
- Element-wise operation
 - memory bandwidth, 329
- ELIZA
 - brittleness, 7
- ELT
 - iteration speed, 133
 - load first, 133

- Embedded GPU accelerator
 - edge deployment, 59
- Embedding
 - definition, 261
 - etymology, 261
 - memory access pattern, 261
 - table, 261
- Embedding Lookup
 - memory-bound operation, 573
- Embedding pruning, 432
- Embedding table
 - capacity bottleneck, 356
 - memory capacity bound, 45
- Emergent behavior
 - definition, 103
- Encoder-decoder
 - definition, 269
- Encryption
 - data lifecycle protection, 833
- End-to-end benchmark
 - complete workflow
 - measurement, 625
- Energy
 - budget, mobile inference, 720
 - data movement cost, 536
 - training scale, 181
- Energy consequence
 - definition, 563
- Energy consumption
 - hardware platforms, 201
 - training costs, 826
- Energy dividend
 - quantization efficiency, 487
- Energy efficiency
 - TFLOPS per watt, 609
 - training cost accounting, 643
- Energy harvesting
 - milliwatt threshold, 64
- Energy hierarchy
 - reference sources, 888
- Energy of transmission
 - local vs. cloud, 44
- Energy per inference
 - binding constraint, 45
 - paradigm comparison, 65
- Energy tax
 - definition, 19
- Energy wall
 - battery constraints, 44
- Energy-movement invariant, 111
- data locality, 846
- decode phase, 733
- distributed training, 408
- Ensemble learning
 - accuracy vs. deployment
 - trade-off, 93
 - deployment cost, 93
- Entanglement
 - ML system coupling, 748
- Entropy
 - appendix refresher, 871
 - definition, 871
- Environment parity
 - production deployment, 741
- Environmental impact
 - carbon accounting, 830
 - ML systems, 826
- Epoch
 - dataset iteration, 211
 - training cycle, 373
 - training cycles, 211
- Equal opportunity
 - definition, 818
- Equalized odds
 - definition, 819
 - postprocessing, 819
- Equivariance
 - translation, 239
- Error Signal
 - backward propagation, 206
- Ethics and governance
 - definition, 29
- ETL
 - reprocessing cost, 133
- EU AI Act
 - risk classification, 823
- systems requirements, 823
- Evaluation pipeline
 - validation metrics, 373
- Exactly-once semantics
 - streaming, 158
- Excessive data movement
 - definition, 579
- Execution problem
 - definition, 291
- Expected calibration error
 - confidence-accuracy gap, 668
- Expected impact
 - definition, 576
- Expected Model Change
 - active learning query strategy, 434
- Experiment tracking
 - tools and practices, 95
- Expert systems, 8
 - knowledge engineering, 173
 - maintenance cost, 173
- Explainability
 - architecture trade-off, 821
 - definition and purposes, 821
 - post-hoc methods (SHAP, LIME), 822
- Exploding gradient problem, 265
 - training instability, 206
- Exponential unit
 - definition, 547
- External tensor payloads
 - definition, 590
- Eyeriss
 - row-stationary dataflow, 570
- F1 score
 - precision-recall harmonic mean, 632
- Face ID
 - privacy architecture, 63
- Facebook AI Similarity Search
 - selection infrastructure, 445
- Fairness
 - accuracy trade-off, 810
 - governance objective, 785
 - individual vs. group, 810
 - mathematical definitions, 810
- Fairness metrics
 - demographic parity, 818
 - equal opportunity, 818
 - equalized odds, 819
 - impossibility theorem, 818
 - loan approval case study, 818
- FAISS, *see* Facebook AI Similarity Search
- Fallacy of peak performance
 - GPU utilization gap, 621
- False Accepts per Hour, 115
- False negative rate
 - fairness impact, 819
- False positive rate
 - always-on systems, 115
- FarmBeats
 - connectivity constraint, 27
- Fault tolerance
 - training checkpoint strategy, 643, 853
- Feature attribution
 - debugging techniques, 782
- Feature Computation
 - hybrid patterns, 135
 - pipeline vs. loader, 135
- Feature computation
 - placement, 135
- Feature consistency
 - training-serving alignment, 741
- Feature contract
 - schema specification, 786
- Feature coverage
 - deployment variation, 671
- Feature Engineering
 - domain knowledge, 138
- Feature engineering
 - classical ML, 173
 - definition, 138
- Feature engineering bottleneck
 - definition, 9
- Feature extraction
 - hierarchical, 242
- Feature importance
 - prediction attribution, 822
- Feature learning
 - hierarchical, 171
- Feature map, 238
- Feature store
 - consistency, 138
 - consistency guarantee, 138
 - definition, 157
 - offline vs. online stores, 158
 - point-in-time retrieval, 156
 - scale, 753
 - streaming updates, 158
 - training-serving bridge, 157
 - training-serving consistency, 753
 - Uber Michelangelo origin, 753
- Federated averaging
 - data minimization, 834
- Federated learning
 - data minimization, 834
 - privacy preservation, 833
- Feedback Loop
 - multi-scale temporal structure, 102
- Feedback loop
 - problem amplification, 825
 - recommendation amplification, 808
 - self-reinforcing bias, 750
 - YouTube case study, 751
- Feedforward network
 - architecture, 192
- Feeding problem
 - definition, 110
- Feeding tax
 - definition, 110
- Field-programmable gate array, *see* FPGA
- FixMatch
 - label-compute trade-off, 437
- FlashAttention
 - 2, improved parallelism, 396
 - 3, FP8 tensor cores, 396
 - appendix refresher, 879
 - framework fusion example, 292
 - I/O-aware algorithm, 393
 - I/O-aware design, 372
 - IO complexity reduction, 395
 - IO-aware, 279
 - memory bandwidth
 - optimization, 393
 - memory optimization, 574
 - memory reduction, 279
 - origin and engineering insight, 393
 - SRAM tiling, 508
 - tiled attention implementation, 596
 - tiling strategy, 394
- FlatBuffers
 - etymology, 694
 - zero-copy serialization, 694
- Flaw of averages, 817
- Flax
 - JAX neural network library, 331
- Fleiss' kappa
 - inter-annotator agreement, 145
- Floating point
 - FP32 precision, 214
- Floating-point operations per second, *see* FLOP/s
- Floating-point unit
 - history, 532
 - precision bottleneck, 532
- FLOP/s
 - peak vs. achieved, 622
- Flow rate
 - definition, 110
- Food and Drug Administration
 - SaMD constraint, 98
- Forced alignment
 - automated labeling, 147
 - speech labeling, 147
- Forgetting events

- coreset selection, 429
- pruning trade-off, 429
- Forward pass
 - compute operations, 377
 - kernel dispatch, 344
 - memory management, 377
- Forward propagation
 - computation process, 197
 - layer computation, 358
 - matrix operations, 198
- Foundation model
 - compute demand, 28
 - definition, 439
 - homogenization risk, 439
- Four pillars framework
 - data engineering, 109
- Four-fifths rule
 - adverse impact threshold, 811
- FP16
 - appendix refresher, 897
 - dynamic range limitation, 388, 389
 - subnormal flushing, 392
 - training precision, 536
- FP32
 - gradient computation, 536
 - numerical precision, 215
- FP32 master weights
 - mixed-precision training, 353
 - optimizer stability, 390
- FP8
 - E4M3 and E5M2 formats, 392
 - hardware acceleration, 392
- FPGA
 - reconfigurability trade-off, 602
- Framework profiler
 - performance analysis tool, 627
 - timeline analysis, 730
- Framework selection
 - constrained optimization, 340
 - trade-off analysis, 340
- Fraud Detection
 - ML vs. rule-based, 83
- Frequency boosting
 - warm-up measurement artifact, 626
- Freshness debt
 - distribution divergence, 159
- Frontier training
 - definition, 26
- Fully-connected layer, 233
 - dense connectivity, 190
- Function unit
 - definition, 547
- Fuse operations
 - definition, 598
- Fused execution
 - definition, 590
- Gabor filters
 - learned equivalents, 174
- GAN
 - synthetic data generation, 440
- Garbage collection pause
 - latency variability, 677
- Gather
 - definition, 542
- Gating mechanism, 263
 - sequence models, 269
- GDDR
 - off-chip DRAM variant, 567
- GDPR
 - Article 22 automated decisions, 823
 - Article 22 requirements, 824
 - compliance, cloud deployment, 53
 - data collection requirements, 122
 - data minimization principle, 834
 - ML erasure constraint, 53
- GELU, 360
- GEMM
 - appendix refresher, 874
 - arithmetic intensity, 328, 891
 - hardware utilization, 290
 - matrix multiplication, 328
 - ML performance predictor, 596
- Gender classification
 - demographic disparities, 809
- Gender Shades
 - algorithmic audit methodology, 811
 - demographic fairness study, 667
 - disaggregated evaluation, 811
- Generalization
 - vs. overfitting, 212
- Generalization gap
 - definition, 28
- Generative adversarial network, *see* GAN
- Generative AI
 - benchmark evolution, 665
- GLUE
 - benchmark saturation, 631
 - language understanding benchmark, 631
- Glue code
 - integration overhead, 750
- Goodhart's law
 - etymology, 616
 - metric optimization, 807
 - metric trap, 616
 - proxy decoupling, 808
- Google Flu Trends
 - proxy drift, 5
- Google TPuv4
 - definition, 557
- GPipe
 - pipeline parallelism, 409
- GPT-2, 228
 - arithmetic intensity, 230
 - autoregressive bottleneck, 45
 - bandwidth-bound decode, 20
 - data selection, 424
 - memory bandwidth, 844
 - training lighthouse, 356
- GPT-3, 830
 - training compute, 356
 - training cost, 640
 - training energy, 11
 - training scale, 10, 50
 - weight storage, 152
- GPU, 826
 - batching utilization, 711
 - etymology, 4
 - history, 533
 - memory hierarchy, 772
 - parallel computation, 201
 - training, power budget, 213
 - utilization monitoring, 771
- GPU backend
 - CUDA implementation, 346
- GPU boost clock
 - benchmark variability, 644
- GPU Interconnect
 - switching fabric, 568
- GPU memory
 - training constraints, 378
- GPU profiler
 - kernel-level analysis, 383
 - system-level analysis, 383
 - timeline analysis, 730
- GPU utilization
 - batch size dependency, 711
- GPU vs. CPU
 - economics, 731
- Graceful degradation
 - data pipelines, 129
 - overload handling, 706
 - triggers, 851
- Gradient
 - backpropagation, 204
 - bias gradients, 206
 - chain rule, 205
 - error signal flow, 206
 - flow in residual networks, 266
 - partial derivatives, 206
 - vanishing and exploding, 206
 - weight gradients, 206
- Gradient accumulation
 - definition, 207
 - framework behavior, 367
 - memory efficiency, 398
 - memory trade-off, 309
 - memory-compute trade-off, 398
- Gradient checkpointing
 - appendix refresher, 879
 - memory reduction, 368
- Gradient compression
 - definition, 397
- Gradient descent
 - definition, 210
 - etymology, 362
 - full-batch computation, 362
 - gradient noise, 363
 - parameter optimization, 362
 - training foundation, 362
- Gradient flow
 - tensor detachment, 312
- Gradient function node
 - computation graph, 311
- Gradient hook
 - inspection and modification, 312
- Gradient inspection, 312
- Gradient instabilities
 - training failure, 171
 - training failure mode, 171
- Gradient-based pruning
 - definition, 472
- GraNd
 - gradient-based coreset scoring, 429
- Graph break
 - causes and detection, 297
- Graph Compilation
 - JIT overhead, 708
- Graph compilation
 - eager-mode capture, 299
 - optimization, 728
- Graph Execution
 - optimization advantages, 292
- Graph Neural Network
 - irregular computation, 579
- Graph optimization
 - computation graph restructuring, 598
 - definition, 597
- Graph retention
 - multiple backward passes, 312
- Graphcore
 - IPU, 550
- Graphics processing unit, *see* GPU
- Great Expectations
 - data validation, 126
- Greedy decoding
 - LLM generation, 723
- Green AI
 - compute reporting, 826
 - efficiency as metric, 826
- Green500
 - energy efficiency ranking, 618
 - HPC energy efficiency ranking, 617
- Ground truth
 - definition, 143
 - proxy limitation, 143
- Group-equivariant CNN, 241
- gRPC
 - inference protocol, 695
- GRU
 - efficiency, 269
 - parameter efficiency, 269
- GTX 580
 - AlexNet training, 9
- Guardrail metric
 - A/B test constraint, 766
- Gustafson's law
 - appendix refresher, 892
- Gustafson, John, *see also* Gustafson's law
- Gzip
 - decompression speed, 155
- H100, 391
- Half-life
 - model decay, 760
- Hardware abstraction
 - framework design, 317
 - two dimensions (data, execution), 318

- Hardware abstraction layer
 - portability, 611
- Hardware acceleration
 - definition, 528
- Hardware accelerator, *see* Accelerator
- Hardware balance, 560
- Hardware cost
 - deployment spectrum, 47
- Hardware lottery
 - architectural bias, 663
 - benchmark bias, 663
- Hardware Mapping
 - combinatorial search space, 580
 - definition, 578
- Hardware mapping
 - hybrid, layer-specific strategy, 596
- Hardware selection
 - workload matching, 610
- Hardware specialization
 - evolution, 531
- Hardware spectrum
 - deployment platforms, 50
 - resource progression, 50
- Hardware sustainability
 - carbon footprint, 609
- Hardware-Software Co-design
 - definition, 529
- Hardware-software co-design
 - edge deployment, 850
- HBM
 - 3D die stacking, 566
 - bandwidth advantage, 276
 - bandwidth-cost trade-off, 539
 - capacity constraint, 344
 - definition, 276
 - memory bandwidth, 394, 734
 - throughput requirements, 539
- Headroom
 - tail latency, 686
- Healthcare AI
 - deployment gap, 87
- Hedging
 - tail tolerance, 705
- HELM
 - multi-dimensional evaluation, 669
- Herding
 - coreset selection, 429
- Heterogeneous computing
 - mobile SoC, 607
- Heuristic
 - kernel selection trade-off, 600
 - pruning, 468
- Hidden layer
 - definition, 207
- Hidden state
 - definition, 247
- Hierarchical processing
 - data flow, 72
- Hierarchical representation
 - depth advantage, 183
- Hierarchical Representation Learning, 281
- Hierarchical selection
 - distributed coreset, 451
- High bandwidth memory, *see* HBM
- High memory latency
 - definition, 580
- High off-chip bandwidth demand
 - definition, 580
- High-risk AI
 - classification criteria, 823
 - EU classification, 813
- Hinton, Geoffrey
 - backpropagation, 180
 - Turing Award, 171
- HIPAA
 - audit requirements, 834
 - compliance, cloud deployment, 53
 - ML system constraints, 824
 - security safeguards, 53
- Hochreiter, Sepp, 269
- HOG
 - fixed vs. learned features, 173
- Holdout test set
 - performance evaluation, 762
- Host-Accelerator Communication
 - transfer bottleneck, 568
- Huang's law
 - GPU performance scaling, 534, 535
- Human-in-the-loop
 - confidence-based routing, 100
 - decision oversight, 815
- Hybrid architecture
 - combining paradigms, 71
 - voice assistant example, 78
 - wake-word detection, 54
- Hybrid machine learning, *see* Hybrid ML
- Hybrid ML
 - definition, 72
 - hierarchical processing, 71
 - integration strategies, 71
 - progressive deployment, 71
 - train-serve split, 71
- Hybrid tiling
 - definition, 594
- Hyperparameter
 - learning rate, 210
 - search cost, 93
 - tuning, 212
- Hypertension
 - ClinAIOPS case study, 796
- i.i.d. assumption
 - held-out evaluation limitation, 672
- I/O bottleneck
 - data loading, 130
- I/O-aware design
 - algorithm optimization, 397
- IB, *see* InfiniBand
- Idempotency
 - definition, 138
 - ML pipelines, 756
 - retry safety, 139
- IIoT
 - latency constraint, 56
- Illustrative latency
 - definition, 547
- im2col, 242
 - convolution to matrix, 543
 - data center vs. mobile, 271
 - memory trade-off, 271
 - memory-compute trade-off, 543
- Image cropping
 - racial bias, 810
- ImageNet
 - benchmark dataset, 118
 - competition progress, 175
 - data scale, 118
 - ILSVRC challenge progression, 667
 - macro benchmark reference, 628
 - scale, 21, 237
 - standardized dataset, 631
 - systems consequence, 237
- Immutable artifact
 - model registration, 755
- Immutable data
 - JAX, 336
- Impact on execution
 - definition, 579
- Imperative programming
 - framework heritage, 335
- Implementation strategy
 - definition, 547
- Importance sampling
 - gradient variance, 458
- In-memory database
 - feature serving, 156
- Incident response
 - ML system failures, 824
 - ML-specific challenges, 780
 - structured processes, 781
- Inductive bias
 - definition, 226
 - etymology, 226
 - hierarchy, 281
 - systems consequence, 226
- Industry 4.0
 - control loop latency, 60
- Inference
 - batch size selection, 215
 - compiler auto-tuning, 647
 - compiler, graph optimization framework, 647
 - deployment, 16
 - deployment phase, 213
 - engine, specialized, 726
 - etymology, 213
 - forward pass only, 202, 213
 - latency optimization, 48
 - latency requirements, 218
 - optimization techniques, 215
 - pipeline phase, 696
 - real-time, storage requirements, 156
 - serverless trade-off, 766
 - zero-copy, mobile optimization, 691
- Inference attack
 - privacy risk, 28
- Inference inflation
 - definition, 703
- Inference runtime
 - execution engine, 725
- Inference server
 - architecture, 693
 - serving architecture, 693
- Inference Serving
 - production infrastructure, 768
- Inference-time compute
 - reasoning models, 690
- InfiniBand
 - bandwidth tiers, 569
 - interconnect performance, 357
 - RDMA for ML training, 569
- Infinity Fabric
 - AMD interconnect, 568
- Information density
 - definition, 871
- Information roofline
 - definition, 866
- Information theory, 870
- Information-compute ratio
 - accuracy-cost frontier, 425
 - definition, 425
 - frontier, 426
- Infrastructure as code
 - ML systems, 770
- Initialization
 - glorot, variance scaling, 196
 - He initialization, 196
 - lazy, deferred compilation, 709
 - Xavier/Glorot, 196
- Input representation
 - definition, 597
- Input-stationary
 - dataflow strategy, 555, 585
- INT8
 - appendix refresher, 897
 - calibration dependency, 651
- INT8 quantization
 - bandwidth trade-off, 846
- Integration bottleneck
 - data movement cost, 538
- Intel
 - Sapphire Rapids, 557
- Intel 8087
 - floating-point coprocessor, 532
 - specialization pattern, 532
- Intel Gaudi 2
 - definition, 558
- Intensity regime
 - definition, 576
- Intentional data engineering
 - fairness evaluation, 815
- Inter-annotator agreement, 144
- Inter-token latency
 - throughput trade-off, 670
- Interaction deduplication, 432
- Interface dependency
 - ML system challenge, 749
- Intermediate representation
 - compiler pattern, 297

- Intermittent computing
 - power management, 65
- Interpretability
 - activation visualization, 220
- Interpretable model
 - transparency by design, 822
- Interrupt shielding
 - latency isolation, 693
- Intersectional analysis
 - combined attribute testing, 817
- Invariance testing
 - fairness verification, 811
- IOPS
 - storage performance, 149
- IoT
 - data wall, 57
 - deployment scale, 60
 - power budget constraints, 826
 - power-constrained environments, 654
- Iron law of ML systems
 - appendix refresher, 861, 892
 - definition, 18
 - deployment analysis, 43
 - equation, 18
 - hardware acceleration, 530
 - multiplicative savings from data selection, 425
 - production monitoring, 775
 - sequential execution, 861
 - synthesis, 842
 - total operations reduction via data selection, 422
 - workflow mapping, 85
- Iron law of training performance
 - definition, 354
 - utilization gap, 354
- Irregular computation patterns
 - definition, 579
- Irregular memory access patterns
 - definition, 580
- Irregular ML components
 - definition, 563
- Iteration tax
 - definition, 83
- Iterative pruning
 - definition, 469
- Jaccard similarity
 - deduplication, 431
- Jacobian matrix, 268
- JAX
 - automatic vectorization, 336
 - composable program transformations, 336
 - composable transformations, 317
 - device-parallel mapping, 336
 - functional transformations, 336
 - gradient transformation, 336
 - JIT compilation, 336
 - pure functions requirement, 336
 - transform-based AD, 317
 - XLA compilation, 337
- jax.grad
 - composable transformation, 336
- Jensen-Shannon divergence
 - distribution shift, 777
- Jetson AGX Orin
 - edge inference, 214
- Jevons paradox
 - inference demand, 687
- JIT compilation
 - fidelity vs. generality, 297
 - runtime optimization, 729
 - scripting, 298
 - shape specialization, 338
 - tracing, 297
- JSON
 - serialization overhead, 695
 - storage overhead, 135
- Jupyter Notebook
 - production risk, 757
- Just-in-time compilation, 301
- k-anonymity
 - privacy technique, 833
- k-center
 - coreset selection, 429
 - coverage guarantee, 429
- Karpathy, Andrej
 - Software 2.0, 4
- Kendall notation
 - queuing models, 703
- Kernel auto-tuning
 - algorithm selection, 726
- Kernel dispatch
 - BLAS kernel selection, 344
 - hardware selection, 328
- Kernel Fusion
 - hardware data movement, 588
- Kernel fusion, 310
 - bandwidth reduction, 310
 - definition, 292
 - eliminating intermediate allocations, 310
 - layer fusion, 726
 - memory efficiency, 589
 - memory traffic reduction, 588
 - optimizing memory-bound ops, 292
 - reducing memory writes, 583
 - static graph optimization, 295, 296
- Kernel launch
 - fixed overhead, 711
 - serving overhead, 711
- Kernel Launch Latency
 - serial overhead, 529
- Kernel profiler
 - hardware-level analysis, 627
- Kernel scheduling
 - hardware utilization, 605
- Kernel selection
 - definition, 597
 - hardware-specific implementation, 599
- Key characteristic
 - definition, 544
- Key management
 - encryption infrastructure, 833
- Keyword spotting, 229
 - automated labeling, 147
 - definition, 21
 - energy constraint, 85
 - sub-megabyte models, 844
 - TinyML archetype, 46
- Killer microseconds
 - latency gap, 698
- Knowledge acquisition bottleneck
 - definition, 8
 - human throughput, 9
- Knowledge distillation
 - benchmark evaluation, 636
 - definition, 477
 - efficiency gains, 826
 - etymology, 476
 - information density, 442
 - student model, 477
 - teacher model, 477
 - teacher-generated targets, 441
 - temperature parameter, 442, 477
- Kolmogorov-Smirnov test
 - continuous drift, 774
 - covariate shift detection, 672
 - drift detection, 101, 125
 - feature distribution, 755
- Kubernetes
 - ML orchestration, 770
- Kullback, Solomon, *see also* Kullback-Leibler divergence
- Kullback-Leibler divergence
 - appendix refresher, 870
 - definition, 128
 - distillation loss, 478
 - distribution comparison, 129, 777
- Kung, H.T.
 - systolic array inventor, 552
- KV cache
 - definition, 260
 - linear growth, 260
 - LLM memory, 724
- LLM serving bottleneck, 772
 - memory capacity, 734
 - memory scaling, 260, 724
 - multi-turn conversations, 719
 - offloading, 724
 - sequence length scaling, 724
 - session reuse, 719
 - transformer memory pressure, 595
- l-diversity
 - privacy technique, 833
- L0 norm
 - parameter count, 468
- Lab-to-clinic gap
 - monitoring requirement, 101
- Lab-to-field gap
 - data distribution shift, 91
- Label propagation
 - graph-based semi-supervised, 436
- Label quality
 - consensus mechanisms, 144
 - systematic checks, 144
- Label quality drift
 - detection, 128
 - silent corruption, 128
- Label shift
 - definition, 128
- Label Studio
 - data quality, 753
- Labelbox, 145
- Laboratory-to-deployment gap
 - performance misalignment, 662
- Ladder of abstraction, 290
- LAPACK
 - framework foundation, 290
 - higher-level numerical routines, 290
- Large language model, *see* LLM
- Late correction cost
 - definition, 83
- Latency
 - bandwidth trade-off, 895
 - decision thresholds, 47
 - deployment constraint, 42
 - deterministic, 693
 - edge deployment, 26
 - etymology, 26
 - first-token, LLM responsiveness, 670
 - inference real-time requirement, 647
 - inference vs. user-perceived, 734
 - mean, monitoring pitfall, 735
 - p50, median response time, 696
 - p95, percentile target, 696
 - percentile, p50 p95 p99 reporting, 632
 - percentiles (p50, p95, p99), 696
 - real-time constraints, 218
 - responsiveness constraint, 42
 - selection, dynamic selection bottleneck, 444
 - serving constraint, 686
 - throughput trade-off, 735
 - vs. throughput trade-off, 28
- Latency budget
 - definition, 696
 - end-to-end analysis, 768
 - optimization objectives, 696
 - request lifecycle breakdown, 696
 - tail latency constraint, 846
- Latency jitter
 - resource contention, 692
- Latency vs. throughput
 - accelerator design, 537
 - paradigm trade-offs, 67
 - trade-offs, 47
- Layer
 - hidden, 189
 - input, 189
 - organization, 189
 - output, 189
- Layer fusion, *see* Operator fusion
- Layer normalization, 268

- Transformer, 258
- Layer sensitivity
 - precision tolerance, 727
- LayerNorm
 - kernel fusion target, 576
 - memory-bound operation, 573
- Layerwise quantization, *see* Quantization
- Lazy Evaluation
 - data processing, 140
- Learnability, 232
- Learnability gap, 232
- Learning rate
 - batch size coupling, 210
 - cosine decay, 367
 - etymology, 210
 - scheduling, 367
 - selection criteria, 212
 - update magnitude, 210
 - warmup, 367
- Learning rate scheduling
 - definition, 367
- LeCun, Yann, 237
 - commercial deployment, 237
 - LeNet, 237
 - Turing Award, 171
- Leibler, Richard, *see also* Kullback-Leibler divergence
- Leiserson, Charles E.
 - systolic array inventor, 552
- LeNet
 - architecture comparison, 218
- LeNet-5, 264
- Li, Fei-Fei
 - ImageNet creation, 118
- LiDAR
 - data volume, 28
- Light barrier
 - latency floor, 42
 - safety-critical systems, 51
- Lighthouse example
 - definition, 573
- Lighthouse model
 - bottleneck probe, 843
 - definition, 20
 - reference workloads, 41
- LIME
 - local explanations, 822
- Limited interconnect bandwidth
 - definition, 579
- Limited on-chip storage
 - definition, 580
- Lineage tracking
 - model provenance, 755
 - model registry, 773
- Linear scaling rule
 - large batch training, 380
- Linear transformation
 - layer computation, 198
- Linnaïmaa, Seppo
 - automatic differentiation, 180
- LINPACK
 - matrix operations benchmark, 616
 - narrow benchmark, 617
- List price
 - definition, 558
- Little's law
 - appendix refresher, 893
 - capacity planning, 702
 - concurrency calculation, 702
- Little, John, *see also* Little's law
- Llama
 - bandwidth-bound decode, 20
 - data selection, 424
 - KV-cache footprint, 772
 - memory-bound inference, 45
- Llama 2
 - 70-billion-parameter
 - memory-bound workload, 848
- Llama 3
 - 8-billion-parameter serving case study, 732
- LLM, 228, 424, 826
 - KV-cache serving, 772
 - labeling assistance, 146
 - training scale, 50
- LLM performance metrics
 - definition, 723
- LLM serving
 - 8-billion-parameter case study, 732
 - production case study, 732
 - token generation, 722
- Load balancer
 - serving infrastructure, 692
- Load shedding
 - retry coordination, 707
- Locality crossover
 - worked example, 59
- Locality of reference
 - ML workload violation, 84
- Locality-sensitive hashing
 - near-duplicate detection, 431
 - probabilistic bucketing, 431
- Log-sum-exp trick
 - appendix refresher, 871
 - numerical stability, 203
- Logic unit cost
 - activation functions, 188
- Logits
 - appendix refresher, 871
 - classification output, 701
 - inference optimization, 188
 - logit function, 871
 - raw network scores, 187
- Loop Ordering
 - computational efficiency, 580
- Loss Function
 - error measurement, 202
 - optimization target, 202
- Loss function
 - cross-entropy, 202
 - etymology, 197
 - landscape geometry, 197
- Loss landscape
 - nonconvex optimization, 210
- Loss Scaling
 - gradient underflow prevention, 313
- Loss scaling
 - appendix refresher, 897
 - gradient underflow prevention, 389
- Lottery ticket hypothesis
 - compression trade-off, 669
 - etymology, 475
- Low-rank factorization
 - definition, 479
- LPDDR
 - low-power DRAM variant, 567
- LSTM
 - compute overhead, 269
 - gating mechanism, 269
- M/D/1 queue
 - deterministic service, 703
- M/M/1 queue
 - Erlang origin, 703
 - wait time formula, 703
- M/M/c queue
 - multi-server, 704
- MAC
 - definition, 235
 - energy cost, 235
 - multiply-accumulate operation, 179
 - neural computation, 179
- Machine learning
 - definition, 8
- Machine learning compiler
 - definition, 597
- Machine learning framework, *see* ML framework
- Machine learning lifecycle, 80
 - definition, 81
 - iteration cycles, 82
 - storage requirements, 155
 - time allocation across stages, 81
- Machine learning systems
 - definition, 12
- Macro benchmark
 - subsystem evaluation, 625
- Magnitude-based pruning
 - definition, 468
- Manifold hypothesis, 232
 - learnability, 232
- MapReduce
 - data locality, 140
- Masked Language Modeling
 - pretext task, 437
- Masked operations
 - definition, 542
- Materialized view
 - training-serving skew, 126
- Matrix engine
 - definition, 557
- Matrix multiplication
 - dense operations, 377
 - forward propagation, 198
 - GEMM optimization, 199
 - neural network workhorse, 542
 - training operations, 358
- Matrix operation
 - hardware acceleration, 544
 - hierarchical decomposition, 543
 - neural network workloads, 549
- Matrix tiling, 274
- Maturity level
 - evolutionary stage, 789
- Maximum mean discrepancy
 - covariate shift detection, 672
- McCrory, Dave, *see also* Data gravity
- McCulloch-Pitts neuron, 231
- Mean time between failures
 - fleet-scale reliability, 853
- Mean time to failure
 - component reliability, 853
- Measurable property
 - responsibility requirements, 837
- Measurement framework
 - data selection metrics, 454
- Measurement variability
 - sources in ML systems, 619
- Mechanical Turk
 - labeling scale, 121
- Medical AI
 - performance metrics, 94
- Membership inference attack
 - privacy testing, 833
 - privacy validation, 833
- Memory
 - capacity, LLM throughput, 734
 - capacity-bound workload, 262
 - constraints, mobile RAM, 721
 - locking, mlock, 693
 - reuse, activation buffers, 215
 - sharing, GPU multi-model, 710
- Memory Access
 - irregular patterns, 563
- Memory access
 - random, 272
 - sequential, 272
 - strided, 272
- Memory access pattern
 - definition, 571
- Memory Bandwidth
 - architectural trade-offs, 567
 - scaling trends, 559
- Memory bandwidth
 - activation function bottleneck, 362
 - bottleneck, 12
 - bottleneck mitigation, 397
 - bound workload, 261
 - LLM bottleneck, 734
 - LLM latency, 734
 - memory wall, 43
- Memory coalescing
 - GPU optimization, 587
 - tensor layout impact, 587
- Memory fetching
 - definition, 588
- Memory footprint
 - computational requirements, 193
 - training overhead, 205
- Memory fragmentation

- KV cache, 718
- Memory hierarchy
 - appendix refresher, 895
 - domain-specific, 535
 - ML training tiers, 376
 - speed-capacity trade-off, 564
- Memory layout
 - NCHW vs. NHWC, 328
 - row-major vs. column-major, 320
 - stride patterns, 320
- Memory management
 - caching allocators, 316
 - definition, 597
 - device placement, 321
 - inference efficiency, 215
 - VRAM requirements, 378
- Memory mapping
 - on-demand model loading, 709
- Memory planning
 - compiler phase, 600
 - definition, 597
 - tensor allocation, 729
- Memory storage order
 - definition, 588
- Memory usage
 - definition, 594
- Memory wall
 - bandwidth divergence, 42
 - compute-memory gap, 42
 - data movement bottleneck, 177
 - execution strategy impact, 291
 - framework execution, 291
 - impact on GPU utilization, 621
 - neural network impact, 177
- Memory-aware tensor layouts
 - definition, 595
- Memory-bandwidth bound
 - precision impact, 727
- Memory-Bound
 - roofline classification, 571
- Memory-bound
 - activation tensors, 699
 - arithmetic intensity, 177
 - operation classification, 291
- Memory-computation tradeoff
 - training systems, 369
- Metric
 - etymology, 632
 - intransitivity, 632
- MFCC
 - audio features, 116, 138
 - dimensionality reduction, 138
- Micro benchmark
 - component isolation, 625
 - diagnostic purpose, 626
- Micro-benchmark suite
 - operator-level testing, 628
- Microcontroller
 - memory ceiling, 65
 - ML deployment, 66
 - TinyML substrate, 26
- Microservice
 - cloud model serving, 690
 - model serving, 690
- MIG
 - hardware isolation, 710
 - multi-model serving, 710
- MinHash
 - distributed, cross-shard deduplication, 451
 - near-duplicate detection, 431
 - signature compression, 431
- Mini-batch
 - gradient descent, 211
- Mini-batch gradient descent, *see* Stochastic gradient descent
- Mini-batch processing
 - accelerator efficiency, 363
- Mixed-precision
 - layer-wise precision assignment, 636
- Mixed-precision Adam
 - storage model, 413, 904
- Mixed-Precision Computing
 - training and inference, 556
- Mixed-precision training
 - benchmarking comparability, 640
 - definition, 353
 - FP16 and FP32, 640
 - FP16 vs. FP32 trade-offs, 313
 - FP16/BF16, 372, 389
 - loss scaling, 389
 - memory savings, 388
 - trade-off, 353
- Mixture-of-experts, *see* MoE
- MixUp
 - image augmentation, 439
- ML Accelerator
 - definition, 537
- ML compiler
 - compilation pipeline, 598
 - definition, 597
 - vs. traditional compiler, 597
- ML development lifecycle
 - vs. traditional software, 25
- ML failure
 - data dominance, 161
- ML framework
 - compiler analogy, 288
 - compiler for the silicon contract, 289
 - compiler role, 41
 - definition, 289
 - differentiation problem, 288
 - dynamic memory management, 370
 - execution problem, 288
 - hardware abstraction problem, 288
 - historical evolution, 290
 - on-device deployment, 63
 - platform comparison, 334
- ML inference benchmark
 - definition, 645
- ML node
 - definition, 741
- ML system benchmark
 - definition, 622
- ML systems hierarchy
 - four layers, 14
- ML test score
 - production readiness, 787
- ML training benchmark
 - convergence throughput scalability, 639
 - definition, 639
- ML workflow
 - systematic framework, 80
- MLCommons
 - nonprofit benchmark consortium, 653
 - open submission, 653
- MLflow
 - experiment tracking, 766
 - model versioning, 152
- MLOps, 29
 - anti-patterns, organizational failures, 790
 - centralized mlops team, organizational pattern, 790
 - definition, 740, 741
 - DevOps comparison, 741
 - emergence, 740
 - federated, embedded engineers, 790
 - production deployment challenges, 740
- MLOps fallacies
 - automated retraining sufficiency, 798
 - automatic consistency, 799
 - DevOps equivalence, 798
- MLOps pitfalls
 - alert routing, 799
 - monitoring overconfidence, 799
 - one-time deployment, 798
 - organizational neglect, 799
- MLP
 - definition, 231
- GEMM-dominated computation, 596
- MLPerf
 - founding, 617
 - founding and purpose, 617
 - inference benchmarking, 576
 - power, ML energy measurement, 617
- MLPerf client
 - local generative AI workloads, 654
- MLPerf Inference
 - benchmark family evolution, 653
 - MultiStream scenario, 722
 - offline scenario, 722
 - server scenario, 719
 - server to edge evaluation, 618
 - serving scenarios, 719
 - SingleStream scenario, 720
- MLPerf Mobile
 - on-device inference, 654
- MLPerf Power
 - energy efficiency measurement, 618
- MLPerf Tiny
 - embedded AI benchmark, 654
 - microcontroller benchmarks, 618
- MLPerf Training
 - data center multi-node scaling, 618
 - standardized framework, 640
- mmap
 - model loading, 709
- MMLU
 - format limitation, 669
 - multitask language understanding, 669
- MNIST, 232
 - dataset saturation example, 673
 - digit recognition, 191
- Mobile assistant
 - definition, 26
- Mobile machine learning, *see* Mobile ML
- Mobile ML
 - battery constraints, 61
 - computational photography, 63
 - definition, 61
 - energy budget, 62
 - energy constraints, 40
 - health monitoring, 63
 - memory bandwidth limits, 63
 - NPU inference, 63
 - resource constraints, 26
 - thermal envelope, 61
 - voice recognition, 63
- Mobile SoC
 - architecture evolution, 607
- MobileNet, 229
- MobileNetV2, 816
 - definition, 21
 - deployment validation example, 615
 - depthwise separable convolutions, 46
 - efficiency gains, 46
 - efficiency under constraint, 844
 - latency-constrained deployment, 244
 - mobile power reduction, 61
 - workflow constraints, 84
- MoCo
 - queue-based dictionary, 439
- Model audit
 - high-stakes systems, 773
- Model benchmarking
 - definition, 615
 - multi-dimensional evaluation, 666
- Model caching
 - container embedding, 709
- Model card
 - documentation standard, 816
 - standardized documentation, 816
- Model collapse
 - recursive synthetic training, 441

- tail distribution compression, 441
- Model compression
 - accuracy-efficiency trade-off, 465
 - definition, 27, 462, 463
 - efficient model representation, 467
 - iterative workflow, 94
 - mobile vision, 61
 - multi-dimensional benchmarking, 651
 - workflow role, 93
- Model debugging
 - decision tree, 781
 - probabilistic failures, 781
 - root cause analysis, 781
- Model decay
 - definition, 774
- Model deployment
 - packaging and serving, 763
- Model extraction
 - model artifact security, 833
- Model FLOPs utilization
 - appendix refresher, 865
 - definition, 382
 - origin, 382
 - PaLM benchmark, 382
- Model governance
 - compliance objective, 785
 - proactive policies, 785
 - transparency and compliance, 785
- Model loading
 - full loading, 709
- Model Memory Pressure
 - architecture-specific, 569
- Model monitoring
 - drift detection, 772
 - performance tracking, 772
- Model Optimization
 - format conversion, 767
 - framework comparison, 767
- Model Parallelism
 - training strategy, 642
- Model parallelism
 - layer-wise partitioning, 406
 - memory constraints, 406
 - multi-accelerator scaling, 606
 - pipeline bubbles, 356
- Model parameter
 - storage requirements, 152
- Model performance metrics
 - definition, 101
- Model quantization, *see* Quantization
- Model registry
 - centralized artifact management, 766
 - provenance tracking, 156
 - version management, 755
 - version promotion, 755
- Model serialization
 - state dictionary pattern, 332
- Model serving
 - definition, 688
 - paradigm shift, 686
- Model size
 - parameter count, 193
- Model type
 - definition, 571
- Model validation
 - candidate selection, 762
 - definition, 96
 - predeployment assessment, 762
- Model versioning
 - artifact tracking, 755
 - deployment patterns, 156
 - requirements, 152
- Model virtualization
 - lifecycle management, 710
- Model weights
 - data governance obligations, 837
 - definition, 5
- Model-centric AI
 - architectural innovation, 672
- Model-infrastructure interface
 - environment parity, 741
- Model-specific execution needs
 - definition, 579
- Model-specific memory needs
 - definition, 580
- Module abstraction
 - hierarchical composition, 330
 - PyTorch nn.Module, 330
- Module hook
 - forward and backward, 333
- MoE
 - definition, 510
 - dynamic computation paths, 564
- Momentum
 - memory overhead, 364
 - optimization physics, 364
- Monitoring
 - health, replica readiness, 692
 - responsible operations, 824
- Monitoring ROI
 - calculation, 780
- Moore's law
 - ML systems consequence, 534
 - slowdown, 534
 - slowdown impact, 535
 - vs. AI growth, 22
- Moravec's paradox
 - compute implications, 9
- Multi-chip scaling
 - beyond single accelerator, 605
 - communication overhead, 606
- Multi-head attention, 257
- Multi-instance GPU, *see* MIG
- Multi-model serving
 - GPU memory management, 710
- Multilayer perceptron, *see* MLP
- Multiply-Accumulate
 - MAC unit, 552
- MultiStream
 - synchronized sensor fusion, 655
- MYCIN
 - expert system, 8
- NAS
 - bi-level optimization, 482
 - bi-level optimization cost, 482
 - definition, 481
 - hardware-aware, 482
 - search cost, 463
- Naïve execution
 - definition, 590
- NCHW
 - channels-first layout, 583, 587
- Netflix
 - prize, production gap, 93
- Network depth
 - training difficulty, 265
- Network interface card
 - node interconnect, 569
- Network RTT
 - latency budget, 768
- Network topology
 - layer organization, 191
- Neural architecture search, *see* NAS
- Neural Network
 - artificial neuron, 178
 - depth vs. width trade-off, 183
- Neural network
 - capacity vs. cost trade-off, 179
 - feedforward, 192
 - layers, 189
 - perceptron, 180
- Neural network application
 - definition, 542
- Neural processing unit
 - energy efficiency, 63
 - mobile devices, 63
- Neuron
 - McCulloch-Pitts model, 178
- Next-Token Prediction
 - pretext task, 437
- NHWC
 - channels-last layout, 583, 587
- NHWC vs. NCHW
 - performance impact, 588
- NIC, *see* Network interface card
- NN
 - dominant MAC pattern, 540
- nn.Module
 - hierarchical composition, 332
 - parameter discovery, 330
 - parameter management, 330
 - registered buffer, 330
 - train vs. eval modes, 331
- No free lunch theorem, 232
 - architecture selection, 232
- Non-maximum suppression
 - irregular memory access, 607
- Non-uniform memory access, *see* NUMA
- Nonlinearity
 - XOR problem, 192
- Normalization
 - batch, 267
 - feature scaling, 137
 - layer, 268
 - parameter persistence, 137
 - RMSNorm, 267
- Notebook to production
 - handoff challenges, 785
- Novelty Effect
 - A/B testing bias, 766
- NPU
 - efficiency trade-off, 607
 - inference benchmarking considerations, 646
 - mobile SoC, 607
 - neural network acceleration, 637
 - on-device strategy, 549
- Nsight Compute
 - profiling tool, 324
- Nsight Systems
 - profiling tool, 324
- Nucleus sampling
 - top-p, 723
- NUMA
 - inference latency, 729
 - memory locality, 729
- Numerical precision optimization, *see* Quantization
- NumPy
 - framework foundation, 290
 - framework lineage, 291
- NVIDIA
 - A100, 558
 - Ampere architecture, 556
 - GeForce 256, 533
 - H100, 537
 - V100, 558
- NVIDIA A100
 - definition, 557
- NVIDIA Accelerator
 - definition, 572
- NVIDIA GTX 580
 - AlexNet memory constraint, 175
- NVIDIA H100
 - definition, 558
- NVIDIA Jetson
 - edge trade-off, 91
- NVIDIA Nsight
 - systems and compute profiling, 627
- NVIDIA Nsight Systems, 385
- NVLink
 - bandwidth hierarchy, 322
 - definition, 568
 - GPU interconnect, 408, 568
 - gradient synchronization, 408
 - intra-node scaling, 557
 - PCIe wall, 557
 - training interconnect, 357, 408
 - training scaling, 568
- NVMe
 - data pipeline, 153
 - GPU utilization, 151
 - storage bandwidth, 376
- Object detection
 - bounding boxes, 143
- Observability
 - system state inference, 773
- Observable degradation
 - continuous measurement, 743
- Offline capability

- paradigm comparison, 67
- Offline scenario
 - maximum throughput mode, 655
- OLAP, 149
- OLTP, 149
- On-call practice
 - ML systems, 783
- On-call rotation
 - ML-specific requirements, 783
- On-chip memory
 - definition, 565
- One-hot encoding
 - classification labels, 202
 - sparsity scaling, 202
- One-shot pruning
 - definition, 473
- oneDNN
 - Intel optimization library, 600
- Online inference
 - definition, 772
- Online softmax
 - incremental computation, 395
- ONNX
 - cross-framework portability, 340
 - definition, 340
 - framework portability, 340
 - hub-and-spoke interoperability, 340
 - model interchange format, 767
- ONNX Runtime
 - cross-platform serving, 725
- Open division
 - MLPerf submission, 654
- OpenCL
 - cross-processor execution, 609
- OpenVINO
 - Intel inference, 726
- Operating system
 - runtime role, 41
- Operation type
 - definition, 544
- Operational maturity
 - system integration, 788
- Operational mismatch
 - ML vs. traditional software, 740
- Operations and monitoring
 - definition, 29
- Operator fusion
 - appendix refresher, 879
 - definition, 507
 - graph compilation, 728
 - inference optimization, 77
 - kernel optimization, 726
 - latency gap elimination, 627
 - memory-bound optimization, 576
- Operator kernel
 - dispatch hierarchy, 327
 - GPU definition, 292, 711
- Optimization
 - Amdahl's law, 76
 - compile-time, constant folding, 728
 - hardware, architecture-specific, 215
 - inference, per-query cost, 827
 - performance, inference, 215
- Optimization difficulty, 232
- Optimization priorities
 - definition, 597
- Optimization priority
 - definition, 576
- Optimization stack
 - multiplicative effects, 454
- Optimization technique
 - definition, 595
- Optimized for
 - definition, 555
- Optimizer
 - Adam, 209
 - adaptive learning rate, 364
 - momentum-based methods, 364
- Optimizer state
 - appendix refresher, 900
 - memory overhead, 353
- Optum, 809
- ORC
 - columnar format, 155
- Oura Ring
 - wearable ML case study, 793
- Out-of-memory error
 - training constraint, 393
- Outlier Detection
 - data cleaning, 137
- Output-stationary
 - dataflow strategy, 555
 - partial sum accumulation, 585
- Over-the-air update
 - edge deployment, 794
 - model distribution, 770
- Overfitting
 - benchmark, over-optimization to test sets, 664
 - definition, 212
 - training challenge, 212
- Overhead-bound
 - small model training, 345
- Overparameterization
 - compression rationale, 467
 - generalization benefit, 175
- Pacing function
 - linear warmup, 433
- Padding
 - batching overhead, 735
- PagedAttention
 - memory efficiency, 718
 - memory fragmentation solution, 718
- Pandera
 - schema validation, 126
- Papermill
 - notebook parameterization, 757
- Parallel Computing
 - Amdahl's Law, 529
- Parallel fraction
 - definition, 529
- Parallel processing
 - neural networks, 201
- Parameter
 - memory multiplier, 189
 - registration, automatic tracking, 330
 - reuse, hierarchical learning, 183
- Parameter freezing
 - transfer learning, 332
- Parameter tracking
 - transformation lineage, 142
- Parameter update
 - optimizer step, 379
- Pareto Frontier
 - serving trade-offs, 703
- Pareto frontier, 227
- accuracy-efficiency trade-off, 668
- compression trade-off, 467, 668
- fairness-accuracy trade-off, 810
- multi-objective optimization, 846
- throughput-latency, 717
- Parquet
 - columnar format, 154, 155
 - I/O reduction, 140
- Partial derivative
 - gradient computation, 206
- Patient treatment loop, 795
- PCIe
 - bandwidth generations, 569
 - host-device interconnect, 321, 569
 - host-to-device transfers, 557
- Peak FLOPS
 - misleading metric, 611
- Peak FP16
 - definition, 572
- Perceptron
 - computational primitive, 231
 - early neural computation, 7
 - etymology, 231
 - linear separability limit, 180
 - Rosenblatt, 180
- Perceptual hashing
 - image deduplication, 431
- Performance gains
 - definition, 594
- Performance per watt
 - efficiency metric, 609
- Performance-per-data
 - metric definition, 454
- Periodic score refresh
 - active learning, 451
- Peripheral component interconnect
 - express, *see* PCIe
- Perplexity
 - etymology, 669
 - filtering, low-information text removal, 432
 - language model prediction measure, 669
 - serving interpretation, 669
- Personalized model
 - user adapters, 719
- Photoplethysmography
 - data quality challenge, 796
 - wearable sensing, 796
- Physical constraints
 - deployment implications, 40
 - memory signaling, 42
 - permanent boundaries, 77
 - speed of light, 42
 - thermodynamics, 42
- Pickle
 - loading overhead, 729
- Pile, The
 - data diversity, 428
- Pinned memory
 - definition, 321
 - DMA transfer, 322, 710
- PipeDream
 - pipeline parallelism, 356
- Pipeline
 - cascading failures, 112
 - debt, workflow fragmentation, 750
 - failure diagnosis, 161
 - failure modes, 124
 - overlapping, latency hiding, 385
- Pipeline Parallelism
 - training strategy, 642
- Pipeline parallelism
 - micro-batch scheduling, 409
 - micro-batching, 409
- Pipelined execution
 - throughput analysis, 43
- Placement considerations
 - definition, 580
- Platt scaling
 - calibration, 97
- Point-in-time correctness, 157
- time travel, 158
- Poisson distribution
 - traffic modeling, 633
- Poisson process
 - queuing model, 703
 - serving arrivals, 719
- Polysomnography
 - label uncertainty, 793
- Pooling
 - global average, 239
- Population stability index
 - appendix refresher, 870
 - credit risk origin, 777
 - definition, 128
 - drift detection, 101, 129, 774, 777
 - drift quantification, 777
- Portrait Mode
 - pipeline latency, 63
- Positive Predictive Value
 - prevalence dependence, 94
- Post-training quantization
 - calibration, 728
 - definition, 485
- Postmortem documentation
 - incident analysis, 781
- Postprocessing
 - confidence thresholds, 218
 - logits to predictions, 700
 - response formatting, 696
- Power budget
 - edge devices, 826

- Power management
 - mobile AI, 608
- Power measurement
 - evaluation triad completion, 656
- Power usage effectiveness
 - data center cooling metric, 659
- Power wall
 - frequency scaling, 42
 - mobile implications, 62
 - thermal limits, 42
- Pre-annotation
 - workflow, 146
- Precision
 - dynamic, adaptive quality, 728
 - hardware design parameter, 555
 - numerical format trade-offs, 896
 - numerical, inference
 - optimization, 215
 - positive prediction accuracy, 632
 - reduced, architectural trade-off, 547
 - throughput trade-off, 727
- Predictive entropy, 129
- Predictive maintenance
 - edge deployment, 60
 - edge duty cycle, 60
- Prefetching
 - accelerator optimization, 535
 - data loader pipeline, 325
 - double-buffering, 154
- Prefill phase
 - compute-bound, 734
 - parallel processing, 734
- Prefix caching
 - KV cache reuse, 724
- Preprocessing
 - digit segmentation, 218
 - divergence, training vs. serving, 707
 - GPU, accelerated pipelines, 698
 - image normalization, 218
 - latency impact, 696
- Pretext Task
 - self-supervised learning, 437
- Price of fairness
 - utility trade-off, 821
- Price/performance
 - definition, 558
- Primary pressure
 - definition, 563
- Primary workloads
 - definition, 557
- Privacy
 - always-listening devices, 833
 - technical protection methods, 833
 - technique hierarchy, 833
- Probabilistic engineering, 6
- Problem definition
 - ML vs. traditional, 88
 - multi-constraint optimization, 88
- ProcessGroup
 - distributed execution, 327
- Processing element
 - accelerator building block, 539
- Production monitoring
 - continuous benchmarking, 675
- Production-monitoring interface
 - feedback loop, 741
- Progressive deployment
 - model compression, 72
- Projection pushdown, 868
- Prometheus
 - monitoring trade-off, 776
- Protocol Buffers
 - serialization, 694
- Proxy metrics
 - definition, 101
- Proxy model
 - practical workflow, 430
- Proxy variable
 - bias detection, 805
 - bias reconstruction, 836
 - definition, 805
- Pruning
 - carbon reduction, 826
 - channel, 469
 - definition, 462, 468
 - dynamic, 472
 - layer, 469
 - low-information filtering, 432
 - neuron, 469
 - Optimal Brain Damage, 467
 - outlier removal, 432
 - quality, label error detection, 432
 - static, 472
 - structured, 471
 - unstructured, 470
- Pseudo-labeling
 - confirmation bias, 436
 - semi-supervised learning, 436
 - semi-supervised technique, 436
- PUE
 - definition, 47
 - efficiency overhead, 47
 - power usage effectiveness, 830
- Pure function
 - JAX, 336
 - JIT requirement, 337
- Python dispatch
 - overhead, 344
- PyTorch
 - autograd implementation, 317
 - autograd system, 311
 - eager execution default, 335
 - preprocessing parameters, 142
 - profiler, 385
 - research-first design, 335
 - tensor layout, 586
- PyTorch profiler
 - performance analysis, 627
- QKV
 - KV cache, 252
- Quality control
 - edge processing, 60
- Quality debt
 - uncorrected errors, 159
- Quantization
 - activation-aware weight (AWQ), 498
 - affine, 497
 - appendix refresher, 897
 - asymmetric, 496, 897
 - AWQ bandwidth reduction, 498
 - channelwise, 497
 - definition, 463, 485
 - energy impact, 488
 - environmental benefits, 826
 - extreme (INT4, binary), 485
 - groupwise, 497
 - inference precision, 214, 321
 - information-theoretic origin, 485
 - INT8 co-design, 529
 - INT8 compression ratio, 615
 - INT8 energy savings, 658
 - layer sensitivity, 727
 - layerwise, 497
 - reduced precision tolerance, 538
 - scale, 497
 - serving optimization, 767
 - serving precision, 727
 - sub-channelwise, 497
 - symmetric, 496, 897
 - TinyML limits, 76
 - zero-point, 497
- Quantization calibration
 - entropy method, 496
 - max method, 496
 - percentile method, 496
- Quantization granularity, *see* Quantization
- Quantization-aware training
 - accuracy preservation, 502
 - definition, 485
- Queries per second
 - inference throughput metric, 650
- Query key value, *see* Query-key-value
- Query-by-committee
 - active learning strategy, 434
- Query-key-value, 252
- Queue contention
 - production latency overhead, 677
- Queueing Theory
 - serving foundations, 736
- Queueing theory
 - serving systems, 702
 - utilization divergence, 703
- RandAugment
 - hyperparameter simplicity, 440
- Random sampling
 - minority representation cost, 815
- Random seed
 - benchmark protocol requirement, 619
- Ray
 - distributed ML, 766
- RDMA
 - multi-node scaling, 569
- Reasoning model
 - inference-time compute, 690
- Recall
 - positive case detection rate, 632
- Receptive field, 238
 - depth trade-off, 238
- Recommendation system
 - DLRM, 46
 - scalability challenge, 114
- Reconstruction attack
 - federated learning vulnerability, 834
- Recurrent neural network, *see* RNN
- Red AI
 - performance at any cost, 826
- Red teaming
 - stakeholder engagement, 811
- Redis
 - online feature serving, 156
- Reduction
 - definition, 542
- Register
 - fastest memory, 565
- Regular dense kernels
 - definition, 563
- Regularization
 - preventing overfitting, 212
- Regulatory compliance
 - AI governance, 823
 - architecture constraints, 834
 - business value, 813
 - EU AI Act, 823
- Relative capacity
 - definition, 565
- Relative latency
 - definition, 565
- Reliability
 - circuit breakers, 130
 - dead letter queues, 129
 - error handling, 129
 - idempotency, 138
- ReLU, 265
 - arithmetic intensity, 891
 - computational efficiency, 187
 - conditional operation overhead, 545
 - dead neuron, 265
 - dying ReLU problem, 187
 - etymology, 187
 - hardware efficiency, 187, 360
- Remote direct memory access, *see* RDMA
- Reorder computations
 - definition, 598
- Replica
 - tail latency improvement, 704
- Representation capacity, 232
- Reproducibility
 - benchmark requirement, 633
 - cross-platform challenge, 662
 - experiment lineage, 95
 - lineage, 142
 - versioning principle, 742
- Reproducible system artifact
 - definition, 93
- Request pipelining
 - GPU utilization, 700
- Request transformation

- tensor formats, 693
- Residual connection
 - Transformer, 258
- Residual learning, 266
- ResNet
 - human-level performance, 175
 - MLPerf reference model, 667
- ResNet-50, 228
 - activation memory, 379
 - arithmetic intensity, 230
 - batching efficiency, 711, 844
 - cloud vs. mobile, 48
 - compute-bound workload, 530
 - data selection, 424
 - data-loading requirements, 130
 - iron-law workload, 20
 - latency budget, 697
 - module hierarchy, 330
 - roofline analysis, 623
 - systems characteristics, 45
 - training memory footprint, 311
 - weight reuse mapping, 595
 - workflow constraints, 85
- Resource isolation
 - serving, 693
- Resource utilization
 - compute and memory, 642
- Responsibility gap
 - definition, 805
- Responsible AI, 29
- Responsible AI engineering
 - definition, 812
- Responsible engineering
 - checklist methodology, 814
 - safety control loop, 804
- REST
 - HTTP protocol, 695
- Retraining, 17
 - continuous updates, 759
 - decision framework, 759
 - economics, cost-benefit optimization, 760
 - scheduled updates, 759
 - scheduled vs. triggered, 759
 - triggered updates, 759
- Retry logic
 - exponential backoff, 129
- Return on compute
 - definition, 20
- Return on investment
 - amortized, 449
 - data selection framework, 449
 - ML systems, 828
 - ML systems, heuristic baseline, 828
- Reweighting
 - bias mitigation, 820
 - importance sampling, 820
- RFM Analysis
 - feature engineering, 138
- Ridge Point
 - hardware balance, 560
 - hardware comparison, 572
- Ridge point
 - appendix refresher, 891, 899
 - definition, 18
 - precision dependence, 371
- Right to erasure
 - data lineage requirement, 833
- Ring AllReduce
 - bandwidth optimization, 409
- RISC-V
 - AI customization, 541
 - vector extensions, 541
- Risk mitigation
 - responsible engineering value, 813
- RMSNorm, 267
 - efficiency, 267
 - LLM adoption, 267
- RMSprop
 - adaptive step sizes, 364
 - origin and Adam inheritance, 364
- RNN
 - definition, 246
- sequential dependency bottleneck, 250
- Role-based access control
 - data governance, 833
- Rollback
 - delayed, 765
 - immediate, 764
 - ML state mismatch, 101
 - model state management, 765
 - rapid, 764
 - strategies, safety mechanisms, 764
- Roofline Analysis
 - conv2d layer, 574
 - dense layer, 575
 - layer normalization, 576
- Roofline Model
 - bottleneck analysis, 48
 - efficiency measurement, 571
- Roofline model
 - appendix refresher, 890
 - compute vs. memory bound, 347
 - definition, 230
 - etymology and origin, 623
 - hardware diagnosis, 531
 - optimization diagnostic, 571
 - training bottlenecks, 370
- Root/power unit
 - definition, 547
- Rosenblatt, Frank, 231
- Rotting asset curve
 - accuracy decay, 774
- Row-oriented storage
 - CSV, 154
- Rule-based programming
 - limitations, 172
- Rumelhart, David
 - backpropagation, 180
- Safetensors
 - secure model loading, 729
 - zero-copy loading, 730
- Safety
 - human oversight requirements, 815
 - responsible engineering, 804
- Sample complexity, 232
- Sampled softmax
 - compute efficiency, 202
- Sampling
 - temperature, top-k, top-p, 723
- Samuel, Arthur
 - machine learning, 8
- SavedModel
 - TensorFlow export format, 334
- Scalability
 - deep learning, 175
 - distributed training metric, 641
 - paradigm comparison, 67
- Scale AI, 145
- Scaling
 - attention, numerical stability, 257
 - frequency, cubic power relationship, 659
 - hybrid, GPU+CPU, 731
 - infrastructure, voice assistants, 54
 - physical ceiling, 412
- Scaling efficiency
 - strong scaling definition, 642
- Scaling laws
 - appendix refresher, 865
 - data limitations, 77
 - data-compute asymmetry, 422
 - power-law origin, 422
- Scatter
 - definition, 542
- Scheduler
 - inference server, 694
- Schema contract
 - data producers, 160
- Schema debt
 - workarounds, 159
- Schema evolution
 - failures, 129
 - silent failure, 129
- Schema validation
 - data quality, 776
 - ML pipelines, 124
 - structural data checks, 776
- Schema-on-read
 - flexibility, 150
 - governance trade-off, 150
- Schema-on-write
 - traditional databases, 150
- Schmidhuber, Jürgen, 269
- Scope creep
 - model cards, 816
- Scratchpad Memory
 - accelerator design, 539
 - software-managed storage, 566
- Scratchpad memory
 - definition, 276, 565
 - explicit control, 276
 - ML advantage, 538
- Secure Enclave
 - biometric privacy, 63
- Security
 - access control architecture, 832
- Select kernels
 - definition, 598
- Selection dependency
 - distributed training, 450
- Selection engineering
 - definition, 444
- Selection inequality
 - definition, 444
- Selection logic
 - definition, 537
- Selective FP32 computation
 - mixed precision, 390
- Selective gradient computation, 312
- Self-attention, 257
- Self-paced learning
 - adaptive difficulty, 433
- Self-supervised learning
 - definition, 437
 - foundation model paradigm, 439
 - iron law restructuring, 437
- Semantic segmentation
 - pixel-level labels, 143
- Semi-supervised learning
 - consistency plus pseudo-labeling, 436
 - definition, 436
 - labeling efficiency, 145
- Sensitivity
 - medical AI metrics, 93
- Separation of concerns
 - MLOps layers, 743
- Sequence length variability
 - batch waste, 717
- Serial fraction
 - Amdahl bottleneck, 529
- Serialization
 - definition, 868
 - overhead, 694
- Server scenario
 - Poisson-distributed traffic, 655
- Server traffic
 - Poisson process, 719
- Serverless AI
 - cold-start benchmarking, 651
- Service Level Objective
 - latency and error rates, 768
- Serving
 - cloud/data center, 690
 - edge inference, embedded serving, 691
 - mobile/edge, 691
 - production deployment, 686
 - TinyML, 691
- Serving inversion
 - throughput-to-latency shift, 686
- Session affinity
 - sticky sessions, 719
- Severity classification
 - incident prioritization, 781
- Shadow deployment
 - inference cost, 764
 - production validation, 764
- Shallow learning

- definition, 171
- Shannon, Claude
 - information theory, 202, 871
- SHAP
 - feature attribution, 822
 - inference latency, 822
- SHAP value
 - model debugging, 782
- Shapley value
 - feature attribution, 822
- Shift handoff
 - context transfer, 784
- Shuffle buffer
 - approximate random sampling, 446
- SIFT
 - variable output size, 173
- Sigmoid
 - computational cost, 185
 - etymology, 185
- Signal-to-noise engineering, 110
- Signal-to-noise ratio
 - definition, 871
- Silent degradation, 16
- Silent failure
 - detection challenges, 807
 - gradual degradation, 743
- Silicon contract
 - constraint propagation, 842
 - definition, 41
 - framework compilation, 289
 - iron law, 43
 - physical constraints, 42
- SimCLR
 - batch requirements, 439
- SIMD
 - CNN inference, 244
 - CPU inference, 729
 - CPU vectorization, 729
 - data-level parallelism, 541
 - definition, 174
 - edge deployment, 244
 - hardware execution, 547
 - vector parallelism, 215
- SIMT
 - definition, 547
- Single instruction multiple data, *see* SIMD
- Single instruction multiple thread, *see* SIMT
- Single-Machine Processing
 - scalability, 140
- Single-metric evaluation
 - fallacy, 676
- Single-node regime
 - definition, 30
- Single-node stack
 - architecture layers, 41
- Single-user traffic
 - mobile serving, 720
- SingleStream
 - sequential inference scenario, 654
- Singular value decomposition
 - rank-bandwidth trade-off, 479
- Skip connection
 - cross-architecture portability, 263
 - gradient flow, 193, 228
- SLA
 - contractual penalty, 696
 - data freshness, 132
 - definition, 648
- SLI
 - definition, 696
- Slice analysis
 - subpopulation debugging, 782
- SLO
 - capacity planning, 702
 - definition, 648
 - percentile targets, 704
 - vs. SLA, 688
 - vs. SLI vs. SLA, 696
- Snappy
 - decompression speed, 155
- Snorkel
 - weak supervision, 146
- SoC
 - always-on island, 115
- Soft affinity
 - load balancing, 719
- Softmax
 - appendix refresher, 871
 - etymology, 188
 - global normalization, 361
 - memory access pattern, 361
 - memory-bound operation, 573
 - numerical stability, 188
 - probability distribution, 187
- Softmax function
 - quadratic bottleneck, 252
 - systems, 252
- Software 1.0
 - definition, 4
- Software 2.0
 - definition, 4
- Sparse connectivity
 - connection patterns, 193
- Sparse scatter
 - definition, 45
- Sparse Tensor Core
 - 2:4 constraint, 550
- Sparsity
 - activation sparsity, 187
 - definition, 511
 - hardware gap, 511
 - hardware speedup requirement, 550
 - memory access irregularity, 564
 - N:M structured pattern, 550
 - SIMD vectorization, 512
 - structured, 2:4 hardware support, 550
- SPEC CPU
 - real application workloads, 617
- SPEC Power
 - energy efficiency, 617
 - server energy efficiency, 617
- SpecAugment
 - audio augmentation, 121
- Special Function Unit
 - activation functions, 539
 - SFU, 544
- Specificity
 - medical AI metrics, 93
- Spectrogram
 - audio features, 138
 - audio representation, 138
- Speculative decoding
 - draft model verification, 718
 - latency reduction, 718
 - throughput, 725
- Speech recognition
 - forced alignment, 147
- Speed of light check
 - utilization diagnosis, 626
- SQL, *see* Structured query language
- SQuAD
 - reading comprehension dataset, 631
 - saturation, 631
- SRAM
 - on-chip cache, 565
 - on-chip fast memory, 565
- SSD cache
 - model loading, 710
- Stable Diffusion
 - generation cost, 441
- Stacking
 - ensemble method, 93
- Stage interface specification
 - contracts, 83
- Staged rollout
 - canary, 97
 - shadow mode, 97
 - workflow, 97
- Staging validation
 - probabilistic adequacy, 764
- Stakeholder communication
 - business alignment, 786
 - technical translation, 786
- Staleness cost
 - accuracy decay model, 760
- Standardization
 - feature scaling, 137
- Static batching
 - fixed batch size, 713
- Static graph
 - define-then-run model, 295
 - definition, 301
 - optimization benefits, 297
- Static inference
 - precomputed predictions, 689
- Static pruning
 - coreset selection, 429
 - definition, 423
 - pretraining filtration, 428
- Static quantization, *see* Quantization
- Static random access memory, *see* SRAM
- Static single assignment
 - appendix refresher, 879
 - compiler IR, 298
- Stationary item
 - definition, 555
- Statistical decision theory
 - loss function origin, 197
- Statistical drift risk
 - accuracy erosion, 846
 - definition, 774
- Statistical learning
 - definition, 9
- Statistical telemetry
 - ML observability, 741
- Statistical validation
 - sampling strategies, 126
- Statistics unit
 - definition, 547
- Stochastic gradient descent
 - etymology, 363
 - gradient estimation, 363
 - memory-utilization trade-off, 363
 - mini-batch procedure, 211
 - Robbins and Monro, 363
 - systems trade-off, 5
- Storage
 - activation, backpropagation, 205
 - architecture, 148
 - cold, archival, 157
 - columnar, ML training, 149
- Store-and-forward architecture
 - intermittent connectivity, 91
- Straight-through estimator
 - gradient approximation, 500
 - quantization-aware training, 500
- Strassen's algorithm
 - matrix multiplication, 358
 - practical crossover, 358
- Stratified evaluation
 - population-based testing, 817
- Stratified sampling
 - fairness evaluation, 815
- Stream ingestion
 - definition, 131
- Streaming multiprocessor
 - GPU architecture, 548
 - occupancy, 548
- Streaming response
 - LLM serving, 724
- Streaming traffic
 - sensor synchronization, 720
- Strong scaling
 - appendix refresher, 892
 - definition, 641
- Structured pruning
 - hardware alignment, 529
- Structured query language
 - feature pipeline primitives, 869
- Structured sparsity
 - N:M pattern, 471, 550
- STUDENT
 - symbolic AI failure mode, 8
- Subgroup imbalance
 - demographic disparity, 671
- Supervised learning
 - labeled examples, 196
- Supported Tensor Core precisions
 - definition, 556
- Sustainability

- environmental responsibility, 826
- Sutton, Richard
 - bitter lesson, 10
- Switch Transformer
 - scaling trade-off, 510
- Symbolic AI
 - definition, 8
- Symmetry-aware architecture, 237
- Synchronization
 - device vs. stream, 324
- Synchronization event
 - fine-grained, 324
 - inter-stream, 324
- Synchronous scoring
 - active learning, 451
- Synthetic data
 - generation, 121
 - rare event coverage, 121
 - scalability, 121
 - simulation, 121
- Synthetic data generation
 - definition, 423
 - three-stage pipeline, 439
- Synthetic Microbenchmark
 - limitations, 620
- System Bottleneck
 - identifying, 43
- System bottleneck
 - dominant constraints, 47
- System composition, 852
- System efficiency
 - feeding tax, 110
 - hardware duty cycle, 700
- System entropy
 - definition, 75
 - model decay, 75, 760
- System is the model, The
 - definition, 842
- System-on-chip
 - energy integration, 63
 - heterogeneous AI acceleration, 607
 - mobile architecture, 63
- Systematic error
 - learned ground truth, 671
- Systems co-design, 9
- Systems thinking
 - principles in ML, 80
- Systolic Array
 - data flow architecture, 552
 - dataflow design, 552
 - energy efficiency, 552
 - TPU configuration, 552
- Systolic array
 - CNN data reuse, 271
 - definition, 535
 - etymology, 271, 552
 - TPU, 271
- t-closeness
 - privacy technique, 833
- T4
 - serving economics, 732
- Tail at scale
 - fan-out amplification, 704
- Tail latency
 - appendix refresher, 869
 - fan-out amplification, 647
 - p99 significance, 647, 849
 - revenue impact, 696
 - utilization threshold, 687
- Tanh
 - etymology, 186
 - zero-centered advantage, 186
- Task scheduling
 - definition, 603
- Task-relevant signal
 - definition, 109
- Technical debt
 - abstraction debt, missing ML interfaces, 750
 - assessment rubric, 787
 - data dependency debt, unstable inputs, 749
 - data engineering, 159
 - dead code, experimental ML paths, 750
 - definition, 746
 - ML compounding, 746
 - ML entanglement, 102
 - ML system complexity, 746
 - ML systems, 102, 825
- Technology S-Curve
 - computing paradigms, 534
- Telemetry
 - ML observability, 741
- Temperature
 - softmax distillation, 477
- Temperature scaling
 - calibration, 97
 - calibration restoration, 669
- Tensor
 - appendix refresher, 875
 - data types (dtypes), 320
 - definition, 318
 - etymology, 318
 - hardware contract, 318
 - memory layout, 171
 - metadata, 319
 - n-dimensional arrays, 171, 318
 - operations, GEMM, 328
 - stride patterns, 320
- Tensor arena
 - TinyML memory, 691
- Tensor Core
 - alignment requirements, 727
 - dimension alignment, 549
 - dimension constraints, 235
 - energy efficiency, 549
 - hardware execution, 549
 - matrix multiplication, 539
 - matrix multiply-accumulate, 391
 - mixed precision, 235
 - mixed precision acceleration, 391
 - mixed-precision acceleration, 313
 - precision support, 549
 - throughput mechanism, 549
 - training throughput, 391
- Tensor decomposition
 - definition, 480
- Tensor-Train, 480
- Tensor layout
 - hardware alignment, 586
 - memory-aware, 583
 - NCHW vs. NHWC, 693
- Tensor memory layout
 - NCHW, 693
 - NHWC, 693
- Tensor parallelism
 - intra-layer splitting, 409
 - operation splitting, 409
- Tensor Processing Unit, *see* TPU
- TensorFlow
 - full framework, 341
 - graph-first architecture, 334
 - preprocessing parameters, 142
 - profiler, 385
 - symbolic differentiation, 317
 - tensor layout, 586
- TensorFlow Data Validation, 126
- TensorFlow Lite
 - mobile deployment, 341
- TensorFlow Lite Micro
 - fixed memory arena, 305
 - microcontroller deployment, 341
- TensorFlow Serving
 - inference server, 693
 - production deployment, 334
- TensorFlow.js
 - browser deployment, 334
- TensorRT
 - accelerator compilation, 578
 - GPU optimization, 726
 - hardware utilization, 339
 - model optimization, 767
- Ternarization
 - definition, 503
- Terraform
 - infrastructure as code, 770
- Tesla D1
 - autonomous vehicle processor, 550
- Tesler's law
 - conservation of complexity, 845
- tf.data, 141
- tf.function
 - graph compilation, 342
- TF32
 - tensor float 32, 556
- TFRecord
 - definition, 154
- Theano
 - computational graph origin, 291
- Thermal design power
 - mobile constraint, 61
- Thermal management
 - mobile constraints, 608
- Thermal throttling
 - impact on sustained perf, 626
 - intelligent migration, 608
 - mobile serving, 720
 - performance limits, 76
 - sustained workload impact, 644
- Thermal trip point
 - hardware safety, 62
- Thermal wall
 - definition, 61
 - mobile constraints, 62
- Thirteen quantitative principles
 - framework, 844
- Thread pinning
 - CPU affinity, 729
- Threshold adjustment
 - fairness mitigation, 820
- Thresholding
 - confidence cutoff, 701
- Throughput
 - accelerator parallelism, 201
 - definition, 641
 - etymology, 641
 - storage bandwidth, 149
 - training, compute vs. data, 154
 - vs. latency trade-off, 632
- Throughput ceiling, 281
- Tied request
 - latency reduction, 706
- Tiered escalation
 - incident response, 784
- Tiered retention
 - hot/warm/cold, 157
- Tiered storage
 - edge deployment, 156
 - hot warm cold architecture, 91
 - ML I/O bottleneck, 91
- Tiling
 - algorithms, matrix computation, 397
 - attention computation, 394
 - definition, 328
 - DRAM traffic reduction, 591
 - memory optimization, 583
 - memory-efficient strategies, 591
 - principle, hardware-graph bridge, 552
 - spatial, data partitioning, 593
 - temporal, 593
- Time per output token
 - decode phase, 723
 - definition, 723
- Time to first token
 - definition, 723
 - prefill phase, 723
- Time travel
 - feature stores, 158
- Time-multiplexing
 - model swapping, 710
- Time-to-accuracy
 - training benchmark primary metric, 641
- Timeline profiling
 - serving optimization, 730
- Tiny constraint
 - definition, 45
- TinyML, 26
 - always-on sensing, 40
 - bare-metal serving, 691

- cost efficiency, 64
- definition, 65
- device form factor, 66
- ecosystem fragmentation, 67
- energy efficiency, 65
- energy gap, 65
- medical wearables, 67
- memory planning, 305
- micro-runtimes, 305
- model compression, 66
- operational constraints, 769
- precision agriculture, 67
- resource constraints, 66
- training memory constraint, 66
- ubiquitous sensing, 64
- wake-word detection, 67
- Token throughput
 - generation rate, 670
- Tool-first anti-pattern
 - fragmented responsibilities, 790
- Top-k
 - sampling, 723
- TOPS
 - comparability limitation, 653
- TOPS/W
 - primary design objective, 656
- torch.compile
 - hybrid execution, 300
- TorchScript
 - tracing and scripting, 297
- Total cost of data
 - formula, 448
- Total cost of ownership
 - cloud vs. edge, 53
 - definition, 827
 - deployment decisions, 53
 - deployment economics, 54
 - inference dominance, 827
 - operational costs, 828
- ToyADAMOS
 - domain gap, 631
- TPU, 827
 - benchmarking considerations, 622
 - data center integration, 537
 - design trade-off, 182
 - domain-specific design, 275
 - matrix unit design, 359
 - origin, 532
 - performance per watt, 275
 - specialization trade-off, 51
 - systolic array scaling, 552
 - XLA compilation, 346
- Trace-based benchmarking
 - production replay, 620
- Traditional compiler
 - definition, 597
- Traditional runtime
 - definition, 603
- Traffic-adaptive batching
 - window tuning, 721
- Train-serve split
 - cost asymmetry, 72
 - economics, 71
- Training
 - bidirectional computation, 48
 - cloud, economics, 758
 - computational cost, 357
 - computational intensity, 353
 - cost asymmetry, 352
 - data throughput, 48
 - inference vs. training, 352
 - infrastructure evolution, 410
 - iron law of training performance, 354
 - mathematical foundations, 357
 - memory decomposition, 878
 - memory equation, 368
 - memory exhaustion threshold, 411
 - memory management, 369
 - memory pressure, 353
 - multi-GPU configuration, 406
 - scaling strategies, 405
 - scaling threshold, 411
 - step, anatomy, 343
 - supervised learning, 196
 - systems, 352
 - time estimation, 893
 - transformer, interconnect
 - sensitivity, 357
 - wall-clock time limit, 411
 - weight optimization, 210
 - weight update, 379
- Training bottleneck
 - compute-bound vs.
 - memory-bound, 382
 - min-of-three relationship, 375
 - profiling, 382
- Training cost
 - electricity, 900
 - scaling trajectory, 352
- Training cost asymmetry
 - definition, 352
- Training loop
 - iterative learning, 211
- Training memory
 - Adam storage model, 888
- Training optimization
 - systematic framework, 384
 - technique selection, 384
- Training pipeline
 - architecture, 372
 - automated ML workflows, 757
 - data flow, 372
 - staged system workflow, 353
- Training profiler
 - bottleneck analysis, 383
 - timeline visualization, 383
- Training systems
 - definition, 29, 352
- Training vs. Inference
 - accelerator design, 536
- Training vs. inference
 - operational differences, 213
 - resource differences, 213
- Training-Serving Consistency, 136
- Training-serving divide
 - definition, 28
- Training-serving skew
 - accuracy impact, 753, 846
 - causes, 126
 - definition, 17, 707
 - diagnosis and prevention, 754
 - serving accuracy fallacy, 735
 - silent accuracy degradation, 707, 754
- Transfer learning
 - compute efficiency, 93
 - definition, 93
 - labeling alternative, 145
- Transform and lighting
 - hardware acceleration, 533
- Transform tensor layouts
 - definition, 598
- Transformation versioning
 - reproducibility, 142
- Transformer
 - architectural shift, 251
 - definition, 256
 - etymology, 251
 - KV-cache memory pressure, 595
 - training evolution, 356
- Transistor tax
 - activation functions, 188
- Translation equivariance, 237
- Translation equivariance, 237
 - systems trade-off, 237
- Translation invariance, 232
- Transparency
 - governance objective, 785
 - regulatory requirements, 822
- Triton
 - inference server, 693
- True positive rate
 - fairness computation, 818
- Trust
 - competitive differentiation, 814
- Tucker decomposition
 - definition, 480
- Turing
 - definition, 556
- Turing, Alan
 - Imitation Game, 6
- Typical operations
 - definition, 576
- Ubiquitous computing
 - etymology, 64
 - TinyML realization, 64
- Uncertainty sampling
 - active learning query strategy, 434
- Undeclared consumer
 - hidden model dependency, 749
 - Tesla case study, 751
- Unified Memory
 - single address space, 569
- Unified memory
 - mobile SoC, 607
- Uniform quantization
 - definition, 494
- Universal approximation theorem, 188, 231
 - depth vs. width, 189
 - systems consequence, 231
- Unreasonable effectiveness of data
 - systems impact, 9
- Unstructured pruning
 - hardware limitations, 514
- Update rule
 - gradient descent, 210
- Upstream echoing
 - definition, 446
- URLLC
 - edge inference constraint, 638
- Useful work per dollar
 - definition, 15
- Utilization
 - high utilization pitfall, 735
 - latency relationship, 703
- V100, 359
 - batching throughput, 711
- Validation
 - accuracy plateau, 212
- Vanishing gradient problem, 250
 - activation saturation, 186
 - exponential decay, 186
- Vector execution
 - definition, 557
- Vector index
 - embedding-based selection, 445
- Vector Operation
 - hardware acceleration, 541
- Vector operation
 - definition, 542
- Vector Unit
 - element-wise operations, 539
- Vector-scalar broadcast
 - definition, 542
- Vectorization
 - automatic batching, 336
 - columnar processing, 868
- Vendor lock-in
 - cloud deployment, 53
 - portability trade-off, 611
- Verification gap
 - definition, 6
 - MLOps closure, 740
 - statistical testing, 846
- Viola-Jones algorithm
 - compute efficiency, 10
- Virtuous cycle
 - definition, 436
- Vision transformer, 270
 - compute trade-off, 270
- ViT, *see* Vision transformer
- vLLM
 - KV-cache pressure, 772
 - virtual memory analogy, 718
- Voice assistant
 - hybrid architecture, 54
- Voice Match
 - enrollment pipeline, 114
- Volta
 - definition, 556
- Von Neumann bottleneck
 - ML accelerator constraint, 559

- VRAM
 - device memory, 710
- Wafer-scale integration
 - single-die approach, 550
- Wake Vision
 - smart doorbell, 15
- Wake-word detection
 - always-on power budget, 67
 - layered architecture, 54
- Wald, Abraham
 - statistical decision theory, 197
- Warehouse-scale computing
 - training evolution, 410
- Warmup
 - cache population, 708
 - cold start mitigation, 708
- Warmup runs
 - benchmarking methodology, 677
- Warp
 - divergence penalty, 548
 - etymology and hardware constraint, 378
 - execution unit, 548
 - GPU execution unit, 378
- Waterfall model
 - contrast with ML, 83
- Wave quantization
 - accelerator utilization, 377
 - tail effects, 377
- Waymo
 - hybrid deployment, 27
- Weak scaling
 - appendix refresher, 892
- Weak supervision
 - cost trade-off, 146
- programmatic labeling, 146
- Web scraping
 - ML training data, 119
- WebDataset
 - sequential sharding, 154
- Weight loading
 - cold start, 708
 - memory efficiency, 711
- Weight matrix
 - layer connections, 189
- Weight Reuse
 - spatial, 595
- Weight sharing, 236
- Weight size
 - definition, 571
- Weight update
 - optimization, 210
- Weight-stationary
 - CNN optimization, 584
 - dataflow strategy, 555
- Weights
 - learnable parameters, 178
 - matrix organization, 189
- Werbos, Paul
 - backpropagation, 180
- Whetstone
 - etymology, 617
 - synthetic floating-point benchmark, 616
- Wildlife conservation
 - TinyML monitoring, 67
- WILDS
 - distribution shift benchmark, 672
- Williams, Ronald J.
 - backpropagation, 180
- Winograd algorithm
 - convolution, 279
 - precision trade-off, 279
- Workflow
 - lab-deployment gap, 80
- Workflow orchestration
 - pipeline automation, 753
- Working set size, 273
- Workload archetype
 - definition, 45
- Workload distribution
 - heterogeneous processor, 607
- Workload imbalance
 - definition, 579
- Worst-Case Execution Time
 - edge AI, 769
- XLA
 - accelerator compilation, 578
 - compilation speedup, 297
 - definition, 297
 - graph compilation, 297
 - warm-up compilation, 768
- XOR problem
 - depth requirement, 191
 - nonlinearity requirement, 192
- YOLOv8-nano
 - edge inference sizing, 59
- Zillow
 - correction cascade, 751
 - distribution shift failure, 825
- Zillow Offers
 - forecast risk, 75