Scaling

Machine Learning Systems

Vijay
Janapa Reddi

Scaling Machine Learning Systems

From Nodes to Fleets: The Engineering of Distributed Scale

Scaling Machine Learning Systems

From Nodes to Fleets: The Engineering of Distributed Scale

Vijay Janapa Reddi
Harvard University

The MIT Press
Cambridge, Massachusetts
London, England

Table of contents

Dedication **xiii**

Author's Note **xv**

Preface **xvii**
Who This Book Is For . . . xvii
Prerequisites . . . xvii
Why a Fleet-Level Textbook . . . xvii
What This Book Covers . . . xix
Suggested Reading Paths . . . xx
How to Read This Book . . . xxi
Companion Resources . . . xxiv
A Living Textbook . . . xxiv
Using This Book in a Course . . . xxiv
Copyright and Licensing . . . xxiv

Acknowledgments **xxv**
Origins . . . xxv
Collaborators and Colleagues . . . xxv
Support . . . xxv

A Note on AI Assistance **xxvii**

Notation **xxix**
Network Communication Models and Collectives . . . xxxiii
Parallelism Dimensions and Scaling Efficiency . . . xxxv
Cluster Reliability, Fault Tolerance, and Checkpointing . . . xxxvii
Distributed Memory Hierarchy and Serving at Scale . . . xxxix
Privacy, Robustness, Governance, and Distributed Constants . . . xliv
Units, Precision, and Disambiguation . . . xlvii

Main Matter

Chapter 1 Introduction **3**
1.1 The Scale Moment . . . 4
1.2 From Single Node to Fleet . . . 9
1.3 AI Scaling Laws . . . 14
1.4 Constraints of Scale . . . 22
1.5 The C^3 Taxonomy and Foundations of Scale . . . 25
1.6 The Fleet Stack and the Lighthouse Archetypes . . . 32
1.7 The Structure of This Textbook . . . 34
1.8 Fallacies and Pitfalls . . . 35
1.9 Summary . . . 36

Part I: The Fleet

Fleet Principles **41**

Chapter 2 Compute Infrastructure **45**
2.1 The Infrastructure Walls . . . 46
2.2 The Accelerator . . . 47
2.3 The Memory Wall . . . 56
2.4 Sustained Compute and Precision . . . 66
2.5 The Node . . . 68
2.6 The Rack . . . 81
2.7 The Pod . . . 93
2.8 Accelerator Selection for Fleet Workloads . . . 96
2.9 Infrastructure Technologies That Move the Walls . . . 99
2.10 Fallacies and Pitfalls . . . 101
2.11 Summary . . . 103

Chapter 3 Network Fabrics **107**
3.1 The Synchronization Backbone . . . 108
3.2 How ML Networking Inverts Data-Center Assumptions . . . 110
3.3 Level 1 Wire and Link . . . 114
3.4 Level 2 Transport and the Performance Model . . . 117
3.5 Level 3 Switch and Topology . . . 124
3.6 Level 4 Fabric Behavior, Congestion, and Routing . . . 134
3.7 Level 5 Cluster Design and Case Studies . . . 141
3.8 Network Virtualization . . . 142
3.9 Monitoring and Debugging . . . 143
3.10 Network Technologies That Move the Ceiling . . . 145
3.11 Fallacies and Pitfalls . . . 147
3.12 Summary . . . 148

Chapter 4 Data Storage **151**
4.1 The Fuel Line . . . 152
4.2 How ML Workloads Invert Storage Assumptions . . . 153
4.3 The ML Storage Hierarchy . . . 156
4.4 The Data Pipeline Equation . . . 175
4.5 GPUDirect Storage and the CPU Data-Copy Bypass . . . 184
4.6 Checkpoint Storage . . . 187
4.7 Versioned Data and Commit Boundaries . . . 192
4.8 Storage Economics . . . 195
4.9 Storage for Serving . . . 198
4.10 Storage Technologies That Move Tiers . . . 200
4.11 Fallacies and Pitfalls . . . 200
4.12 Summary . . . 202

Part II: Distributed ML

Distributed ML Principles **207**

Chapter 5 Collective Communication **211**
5.1 From Parallelism to Communication Patterns . . . 212
5.2 The Physics of Data Movement . . . 213
5.3 Parallelism Strategies and Their Traffic . . . 215
5.4 Collective Operation Primitives . . . 216
5.5 Pricing Collectives by Rounds and Chunks . . . 220
5.6 AllReduce Algorithms and Topologies . . . 222

5.7 Hierarchical Communication . . . 233
5.8 AllToAll Exchanges at Scale . . . 239
5.9 Gradient Compression Under Bandwidth Scarcity . . . 243
5.10 Communication-Computation Overlap . . . 249
5.11 Communication Libraries and Runtimes . . . 253
5.12 Fallacies and Pitfalls . . . 254
5.13 Summary . . . 256

Chapter 6 Distributed Training **259**
6.1 Why Distribution Is Necessary . . . 260
6.2 The Parallelism Design Space . . . 263
6.3 Data Parallelism . . . 265
6.4 Scaling Efficiency and Convergence . . . 278
6.5 Model Parallelism . . . 291
6.6 Hybrid Parallelism . . . 308
6.7 Multi-Model Training for Alignment and RLHF . . . 313
6.8 Parallelism Strategy Comparison . . . 318
6.9 Frameworks and Runtime Contracts . . . 320
6.10 Fallacies and Pitfalls . . . 321
6.11 Summary . . . 323

Chapter 7 Fault Tolerance **325**
7.1 Failure Analysis at Scale . . . 326
7.2 The Mathematics of Inevitable Failure . . . 327
7.3 Failure Taxonomy . . . 331
7.4 Hardware Faults at Fleet Scale . . . 336
7.5 Check-and-Verify Defenses Against Silent Data Corruption . . . 340
7.6 Software Faults . . . 343
7.7 The Cost of Failure . . . 345
7.8 Checkpointing and Progress Preservation . . . 347
7.9 Failure Detection and Recovery . . . 355
7.10 Elastic Recovery . . . 364
7.11 Serving Fault Tolerance . . . 368
7.12 Graceful Degradation . . . 374
7.13 Distributed Debugging and Observability . . . 377
7.14 Fault Injection Tools and Frameworks . . . 382
7.15 Case Studies . . . 382
7.16 Fallacies and Pitfalls . . . 384
7.17 Summary . . . 386

Chapter 8 Fleet Orchestration **389**
8.1 The Scheduling Problem . . . 390
8.2 Distributed Scheduling Complexity . . . 392
8.3 Scheduling Objectives and Their Conflicts . . . 395
8.4 Core Allocation Constraints . . . 397
8.5 Orchestration Paradigms . . . 404
8.6 Topology-Aware Scheduling . . . 411
8.7 Elastic Scheduling . . . 418
8.8 Cost Optimization . . . 423
8.9 Custom ML Schedulers . . . 429
8.10 Serving Capacity in the Shared Fleet . . . 432
8.11 Multi-Tenancy and Quotas . . . 436
8.12 Debugging Cluster Utilization . . . 440
8.13 Fallacies and Pitfalls . . . 444
8.14 Summary . . . 445

Part III: The Fleet in Production

Fleet-in-Production Principles **451**

Chapter 9 Fleet Efficiency **455**
9.1 Where Fleet Capacity Goes . . . 456
9.2 ML Goodput . . . 457
9.3 Measuring Goodput . . . 465
9.4 Scheduling Goodput . . . 473
9.5 Runtime Goodput . . . 477
9.6 Program Goodput . . . 482
9.7 Recovering Program Goodput . . . 490
9.8 The Whole Fleet . . . 499
9.9 Trading Between Factors . . . 503
9.10 Fallacies and Pitfalls . . . 512
9.11 Summary . . . 515

Chapter 10 Serving at Scale **517**
10.1 Inference Cost and Serving Architecture . . . 518
10.2 Limits of Single-Machine Serving . . . 519
10.3 SLO Goodput and the Serving Tax . . . 521
10.4 Serving Cost Dominance in Production . . . 522
10.5 Serving Workloads and the Serving Hierarchy . . . 523
10.6 Batching Strategies at Scale . . . 526
10.7 Memory and Decode-Time Management . . . 543
10.8 Prefill and Decode . . . 556
10.9 Weight Quantization for Serving . . . 561
10.10 Model Sharding for Inference . . . 566
10.11 Load Balancing and Request Routing . . . 575
10.12 Multi-Tenancy and Isolation . . . 582
10.13 Autoscaling and Capacity Control . . . 585
10.14 Global Inference Infrastructure . . . 589
10.15 Case Studies . . . 595
10.16 Fallacies and Pitfalls . . . 601
10.17 Summary . . . 603

Chapter 11 Edge Intelligence **607**
11.1 The Edge Learning Paradigm . . . 608
11.2 Why Learn on the Device . . . 610
11.3 Design Constraints . . . 613
11.4 Model Adaptation . . . 625
11.5 Data Efficiency . . . 634
11.6 Federated Learning Algorithms . . . 639
11.7 Federated Learning Systems at Scale . . . 644
11.8 Continual Adaptation Under Edge Budgets . . . 655
11.9 Integrating the Edge Intelligence Pillars . . . 656
11.10 Production Integration . . . 659
11.11 Engineering Challenges and Mitigations . . . 664
11.12 Fallacies and Pitfalls . . . 667
11.13 Summary . . . 668

Chapter 12 ML Operations at Scale **671**
12.1 Why a Model Fleet Needs a Control Plane . . . 672
12.2 The N-Models Problem . . . 673
12.3 Governance of State and Change . . . 683
12.4 Feature Store Operations . . . 688

12.5 CI/CD for ML at Scale . . . 697
12.6 Monitoring at Scale . . . 718
12.7 Production Debugging and Incident Response . . . 728
12.8 Platform Engineering . . . 734
12.9 Fleet Economics and Utilization . . . 738
12.10 Organizational Operation and Learning . . . 758
12.11 Case Studies . . . 761
12.12 Fallacies and Pitfalls . . . 762
12.13 Summary . . . 764

Part IV: The Responsible Fleet

Responsible Fleet Principles **769**

Chapter 13 Security & Privacy **773**
13.1 The Expanded Attack Surface . . . 774
13.2 Security and Privacy Foundations . . . 775
13.3 Learning from Security Breaches . . . 778
13.4 Systematic Threat Analysis and Risk Assessment . . . 780
13.5 Model-Specific Attack Vectors . . . 785
13.6 Fleet-Internal Trust Boundaries . . . 803
13.7 Hardware-Level Security Vulnerabilities . . . 806
13.8 Comprehensive Defense Architectures . . . 814
13.9 Differential Privacy . . . 831
13.10 Selecting and Sequencing Defenses . . . 841
13.11 Fallacies and Pitfalls . . . 844
13.12 Summary . . . 845

Chapter 14 Robust AI **847**
14.1 The Silent Failure Problem . . . 848
14.2 Real-World Robustness Failures . . . 850
14.3 A Unified Framework for Robust AI . . . 852
14.4 Environmental Shifts . . . 857
14.5 Adversarial Attacks . . . 869
14.6 Adversarial Defenses . . . 873
14.7 Data Poisoning Defenses . . . 879
14.8 Robustness of Generative Models . . . 884
14.9 System-Level Faults . . . 886
14.10 Fallacies and Pitfalls . . . 888
14.11 Summary . . . 890

Chapter 15 Sustainable AI **893**
15.1 The Energy Ceiling . . . 894
15.2 The Scale of Environmental Impact . . . 896
15.3 Exponential Growth vs. Physical Constraints . . . 897
15.4 Energy Measurement and Modeling . . . 903
15.5 Carbon Footprint Calculation . . . 912
15.6 Data Center Energy and Resource Consumption . . . 923
15.7 Training vs. Inference Energy Analysis . . . 928
15.8 Hardware Lifecycle and E-Waste . . . 937
15.9 Mitigation Strategies . . . 943
15.10 Policy, Regulation, and the Path Forward . . . 953
15.11 Fallacies and Pitfalls . . . 956
15.12 Summary . . . 958

Chapter 16 Responsible AI **961**

16.1 The Governance Imperative . . . 962
16.2 The Governance and Assurance Lifecycle . . . 963
16.3 Core Principles and the ML Lifecycle . . . 966
16.4 Architectural Trade-Offs . . . 967
16.5 Fairness and Bias Mitigation . . . 971
16.6 Privacy and Data Governance . . . 989
16.7 Explainability and Transparency . . . 995
16.8 AI Safety, Robustness, and Generative Controls . . . 1001
16.9 Accountability, Governance, and Sociotechnical Dynamics . . . 1007
16.10 Governing Responsible AI at Fleet Scale . . . 1017
16.11 Fallacies and Pitfalls . . . 1025
16.12 Summary . . . 1027

Synthesis

Chapter 17 Conclusion **1031**
17.1 Distilling Distributed ML Systems . . . 1032
17.2 Six Principles of Distributed ML Systems . . . 1032
17.3 The Complete Production System . . . 1037
17.4 Competencies Mastered . . . 1041
17.5 Where the Binding Constraint Moves Next . . . 1041
17.6 Fallacies and Pitfalls . . . 1042
17.7 Summary . . . 1043

References

References **1047**

Appendices **1075**

Appendix A Single-Machine Foundations **1075**

Appendix B The C^3 Taxonomy **1085**

Appendix C Fleet Foundations **1099**

Appendix D Communication Foundations **1113**

Appendix E Reliability Foundations **1123**

Appendix F Inference Foundations **1135**

Appendix G System Assumptions **1145**

Glossary **1159**

Index **1171**

0

Dedication

For Kari,
who gave me the time this book required
and never asked for it back.

And for Nora and Maya,
the reason any of it matters.

0

Author's Note

When a new model is announced, the world looks at the model. It almost never looks at the fleet. Behind many frontier models, however, stand thousands of accelerators held in lockstep across a network where one slow link can stall the whole run, cooled by systems that move megawatts of heat, kept alive by protocols that treat failure not as an exception but as a scheduled event. None of this fits in the announcement. The model gets the paper. The fleet gets a footnote, if that.

This book is an argument that the fleet deserves more than a footnote. We have built vast distributed systems before, but none quite like this one, because none of the others had to be right about two things at once. The telephone network had to stay available; it never had to converge. The internet had to deliver the packet; it never had to preserve the gradient inside it. A fleet has to do both. It must move data the way any distributed system does, and it must keep the mathematics of learning intact while it moves. A node that drops out does not only cost capacity; it can break the synchronization a training step depends on and force a rollback that wastes work since the most recent checkpoint. That double obligation, computational and statistical, sustained at the scale of a small campus, is the engineering this book makes visible.

The obligation does not end when training does. A model that cannot be trained is an idea, not a system. A model that cannot be served is a research result, not a product. A model that cannot be governed is a liability, not an asset. At every stage, from the first allocated node to the last served request, the fleet decides what is possible and what remains imagined.

The problem under every chapter is easy to state and hard to solve: what it takes to build not one machine learning system but a thousand, make them work as one, and run them responsibly. I believe that is a defining engineering problem of this generation, and that the people who take it on, the ones who rarely get the headline, deserve a book that treats it as one.

— Vijay Janapa Reddi
Cambridge, Massachusetts
2026

0

Preface

Who This Book Is For

This book is for anyone who understands machine learning on a single machine and now faces the challenge of making it work at scale. The problems of scale are not the problems of a single node, only bigger. They are qualitatively different: networks partition, hardware fails as a statistical certainty, and the societal impact of every design decision is amplified by millions of users.

The assumed foundation is what an introductory ML systems course or equivalent experience supplies, with no particular prior textbook required. The advanced scope here is the transition from that single node to a coordinated fleet. That scope ranges from designing the physical substrate of an AI data center, to scaling training beyond the limits of a single accelerator, to reasoning about what happens when a production fleet serves a global user base. Throughout, the focus is on the physics of distribution, the invariants that govern every fleet-scale ML system regardless of the framework, the model, or the era.

Prerequisites

Readers should bring the following background:

- **Single-machine ML systems**: Understanding of ML workflows, neural network training, model optimization, and deployment on a single machine. Readers should be comfortable reasoning about memory hierarchies, computational graphs, accelerator execution, and hardware-software interactions. Appendix A recaps the single-machine performance model (the iron law, arithmetic intensity, and the Roofline Model), the memory hierarchy, training memory per parameter, and transformer compute that later chapters assume.
- **Programming**: Fluency in Python, including functions, classes, and data manipulation with NumPy. Familiarity with at least one ML framework (PyTorch, TensorFlow, or JAX).
- **Mathematics**: Comfort with linear algebra, basic calculus, and probability at the undergraduate level. Section E.1.1 restates the probability results that the failure analysis uses, and section F.1.1 develops the queuing results that serving analysis uses.

The following background is helpful but not required:

- **Distributed systems**: Familiarity with networking concepts (TCP/IP, bandwidth, latency), parallelism, and basic distributed computing deepens understanding of the fleet infrastructure and training chapters.
- **Cloud infrastructure**: Experience with containerization and orchestration (Docker and Kubernetes) or cluster schedulers (Slurm) provides useful operational context.
- **Production engineering**: Understanding of monitoring, observability, and site reliability practices will enrich the deployment and operations chapters.

Why a Fleet-Level Textbook

In 2012, training AlexNet took five to six days on two GPUs (Krizhevsky et al. 2012). By 2023, public GPT-4-class estimates placed frontier training at roughly 25,000 GPUs running for about three months (Patel and Wong 2023); OpenAI did not disclose the actual hardware configuration (OpenAI et al. 2023). This is not the same engineering problem at a larger scale. It is a categorically different engineering problem.

On a single machine, performance is governed by the memory wall, the gap between processor speed and memory bandwidth. In a distributed cluster, a new limit emerges. Bisection bandwidth, the total bandwidth across the narrowest cut that divides the cluster into two halves, caps the throughput of any traffic that must cross that cut. Data no longer moves through local hierarchies alone; it traverses optical fabrics governed by the speed of light between racks and across continents. Hardware failures, rare enough on a single machine to be treated as exceptions, become routine statistical events at fleet scale. Long training runs must plan for interruptions rather than rely on every component remaining available until completion.

These are not operational annoyances. They are physical constraints, as fundamental and as permanent as the memory wall, Amdahl's Law, or the energy cost of data movement. They require their own engineering discipline, built on the invariants this book develops:

- **The Empirical Scaling Law**: Within empirically validated regimes, model loss falls as a power law in training compute, data, or parameters, so each further improvement demands multiplicatively more resources. This is the pressure that moves machine learning from single machines onto fleets.
- **The Bisection Bandwidth Theorem**: For communication patterns that must cross the minimum bisection, throughput is bounded by that bisection bandwidth.
- **The Fleet Law**: Under fixed-work scaling, reduced per-node computation competes with communication and coordination costs, so adding workers does not guarantee proportional speedup.
- **The Young-Daly Checkpoint Interval**: The optimal checkpoint interval balances the cost of writing state against the expected cost of lost work due to failure.
- **Serving Cost Dominance**: For high-traffic production models, cumulative inference operational expenditure can exceed the one-time training cost, making serving efficiency a lifecycle optimization target.
- **The Fairness Impossibility Theorem**: With unequal group base rates and imperfect prediction, risk scores cannot satisfy groupwise calibration and equal mean scores across groups within both outcome classes simultaneously. Thresholded error-rate parity and demographic parity are distinct requirements with their own trade-offs.

If single-machine ML systems engineering asks *"what can one machine build?"*, then fleet-scale ML systems engineering asks *"what can a thousand machines build together, and at what cost to reliability, efficiency, and society?"* That final clause is what sets fleet-scale engineering apart. A single machine affects its user; a fleet affects everyone its outputs touch, across millions of decisions per second. At that scale, reliability, efficiency, and fairness stop being features and become engineering constraints as binding as bandwidth or power. A system that scales but cannot be trusted has not been engineered. It has only been enlarged.

This book develops that discipline. It follows the Hennessy and Patterson pedagogical model, extending the quantitative, principles-first approach to the distributed scale (Hennessy and Patterson 2011; Hennessy and Patterson 2019, 2017). Just as their work taught a generation to reason about processor performance from first principles, this book teaches the reader to reason about fleet performance, replacing intuition about distributed systems with measurement, and ad hoc scaling with engineering.

The architectural priorities and operational failure modes explored in these chapters did not originate in an academic vacuum. They were gathered directly from practitioners in industry, through technical exchanges at conferences, engineering meetups, and conversations with colleagues. These colleagues generously made the time to explain what breaks when scaling models across thousands of accelerators and what practicing systems engineers must understand.

Single-node foundations supply the bricks of those systems: tensors, kernels, memory hierarchies, and optimization loops. What breaks at fleet scale is rarely an individual brick, because a fleet is closer to an orchestra than to a larger wall. An orchestra is not simply many musicians playing at once. It needs a conductor, a score, precise synchronization, and the ability to recover when a player drops out, none of which governs a solo performer. A training cluster, a serving fleet, and an edge deployment are vastly different systems, but beneath them sit the same building blocks: parallelism

strategies, collective communication primitives, checkpoint protocols, scheduling algorithms, and the trade-offs between throughput, latency, fault tolerance, and cost.

What This Book Covers

This book focuses on the machine learning fleet: coordinated machines that jointly train, serve, or maintain ML models. Warehouse-scale training clusters are one form; geographically distributed serving replicas and heterogeneous edge devices are others. Across these forms, communication and coordination couple local execution to fleet behavior, while reliability, resource budgets, and governance constrain the design. The scope is the fleet as it exists today and the governance of operating it at scale. Research-frontier topics such as fully autonomous fleet management, next-generation interconnects, neuromorphic computing at scale, and formal verification of fleet-level safety properties lie beyond it.

The content is organized into four parts:

- **Part I: The Fleet** builds the physical substrate of the AI data center from the ground up: the silicon and cooling of compute infrastructure, the network fabrics that connect thousands of accelerators, and the scalable data storage that feeds training pipelines.
- **Part II: Distributed ML** establishes the logic of distribution beyond a single machine: the collective communication primitives that synchronize gradients, the parallelism strategies (data, tensor, and pipeline) for training models too large for single devices, the fault tolerance mechanisms that ensure reliability when failure is routine, and the orchestration systems that manage the fleet.
- **Part III: The Fleet in Production** takes the trained model from the cluster to the world: fleet efficiency, which measures how much of a fleet's capacity becomes useful work, serving at scale for millions of concurrent requests, edge intelligence across heterogeneous and federated devices, and the operational lifecycle required to manage production fleets.
- **Part IV: The Responsible Fleet** confronts the forces that determine whether the fleet serves its users or harms them: security and privacy, robust system design, environmental sustainability, and responsible engineering as constraints as fundamental as power density or bisection bandwidth.

These four parts trace a path through the fleet stack, a layered abstraction that organizes an ML system at distributed scale, from the physical facility up to its governance. The stack has eight rows grouped into four bands, and each band is one part of the book. At the base, the Infrastructure band (Part I) holds **Facility & Accelerators** and **Fabric & Storage**, the physical substrate. On top of that hardware, the Distribution band (Part II) adds **Distributed Execution** and **Reliability & Control** to manage parallel work, communication, scheduling, and recovery. Once a model can be trained reliably, the Serving band (Part III) provides **Runtime & Serving** and **Operations** to take models to users and manage the production lifecycle. At the top, the Governance band (Part IV) layers **Assurance** (security, privacy, and robustness) and **Governance & Sustainability** to turn trust, risk, policy, and environmental cost into engineering constraints. Each row provides an abstraction to the row above and consumes the row below. Engineering decisions in the Infrastructure band constrain what is possible in the bands above it, while governance requirements constrain choices throughout the stack. Privacy, acceptable risk, and environmental budgets enter the design from the beginning; Part IV develops their enforcement mechanisms after the infrastructure and operational dependencies are available.

Where the fleet stack locates responsibility, a second framework diagnoses cost. The C^3 taxonomy divides the wall-clock time of distributed execution into Compute, Communication, and Coordination, and the Fleet Law turns that division into an equation for distributed step time. Every chapter applies this lens to its own layer of the stack. Three recurring workloads, the archetypes, keep the analysis concrete: a large dense language model (Archetype A), a recommendation model with very large embedding tables (Archetype B), and a model trained by federated learning across phone-class devices (Archetype C). Each stresses a different part of the fleet, and the chapters return to whichever one exposes their central constraint.

Each part builds on the previous one, so sequential reading is recommended. The reading paths below offer alternatives for readers with specific goals.

Suggested Reading Paths

The chapters are ordered so that each can draw on everything before it, but not every chapter consumes the one immediately before it. A dependency map derived from the chapters' own cross-references records which chapters build directly on which (figure P.1). An arrow in the map means the later chapter uses a model, mechanism, or constraint that the earlier one establishes. Two features of the map guide the choice of path. First, every chapter lies on a route to chapter 17, whose six principles cite every chapter of Parts I through IV, so no chapter is a dead end. Second, several dependencies cross part boundaries: distributed training builds on both the storage pipeline of chapter 4 and the collectives of chapter 5, and fleet efficiency (chapter 9) builds on chapter 8. Consecutive chapters without an arrow mark a change of thread rather than a gap. Chapter 5, for example, builds on the fabric model of chapter 3 rather than on the storage chapter that precedes it.

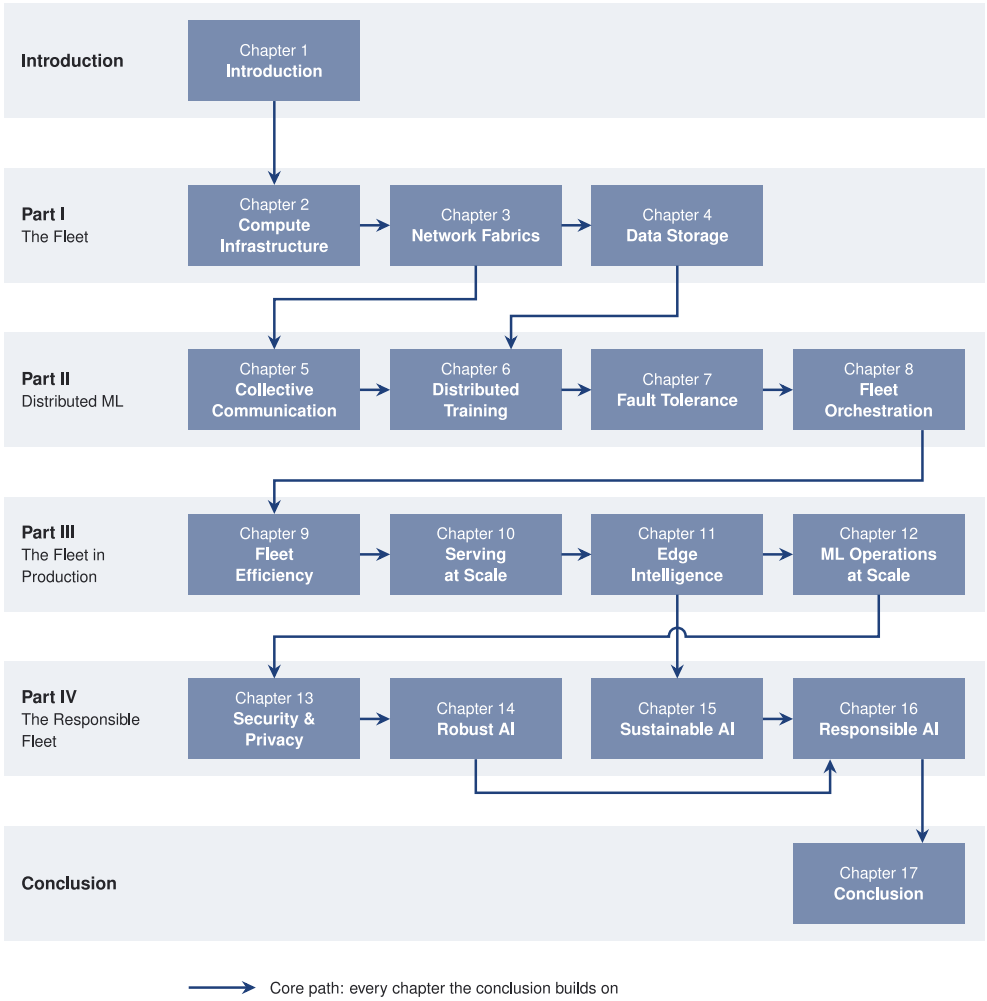

Figure P.1: Chapter Dependency Map: Each arrow points from a chapter to a later chapter that builds on its results, derived from the book's cross-references with transitive links omitted. Every chapter lies on the core path the conclusion builds on, so the map draws no optional branch. Book order is one valid reading order; a missing arrow between consecutive chapters means the later chapter does not directly consume the earlier one's results.

Each path below begins with the introduction, assumes the single-machine prerequisites listed earlier in this preface, and keeps every chapter that its destination depends on in the map:

- **The Complete Path** (all chapters, in sequence): Follows the causal chain of constraints that governs the ML fleet from first principles of distribution. After the introduction, progress through Parts I, II, III, and IV.
- **The Infrastructure Path** (physical fleet and orchestration): Focuses on Part I (chapter 2, chapter 3, chapter 4), then chapter 8. Before orchestration, use the collective traffic patterns in chapter 5, the workload partitioning in chapter 6, and the checkpoint/recovery model in chapter 7 to establish what the scheduler must place, connect, and preserve.
- **The Scaling Path** (distributed training depth): Establishes the memory-capacity, network-bandwidth, and storage-throughput models in Part I before focusing on chapter 5, chapter 6, chapter 7, chapter 8, and chapter 9. Readers already comfortable estimating those limits can use Part I as a reference.
- **The Production Path** (deployment and operations): Uses Part I for fleet capacity and network constraints. It then establishes collective communication in chapter 5, model partitioning in chapter 6, failure and recovery in chapter 7, placement and isolation in chapter 8, and how much of the fleet's capacity becomes useful work in chapter 9. Continue through chapter 10, chapter 11, chapter 12, and Part IV (chapter 13, chapter 14, chapter 15, and chapter 16).

How to Read This Book

Across any of those reading paths, recurring visual cues in the margins orient the reader within the fleet stack. At every chapter opening, a margin figure highlights which rows of the stack the chapter addresses. The full stack below shows each row with its scope, grouped into the four parts of the book:

The Roman numerals (I–IV) mark those four parts, while the row intensity marks emphasis. The stack uses one accent color throughout so that shade, not hue, carries the chapter-focus signal.

In the margin figures throughout the book, the intensity of each row reflects how strongly the chapter engages it. Brighter (higher intensity) marks the chapter's primary focus; dimmer (lower intensity) marks contextual relevance as constraints propagate upward. Consider three chapters that illustrate how the emphasis shifts:

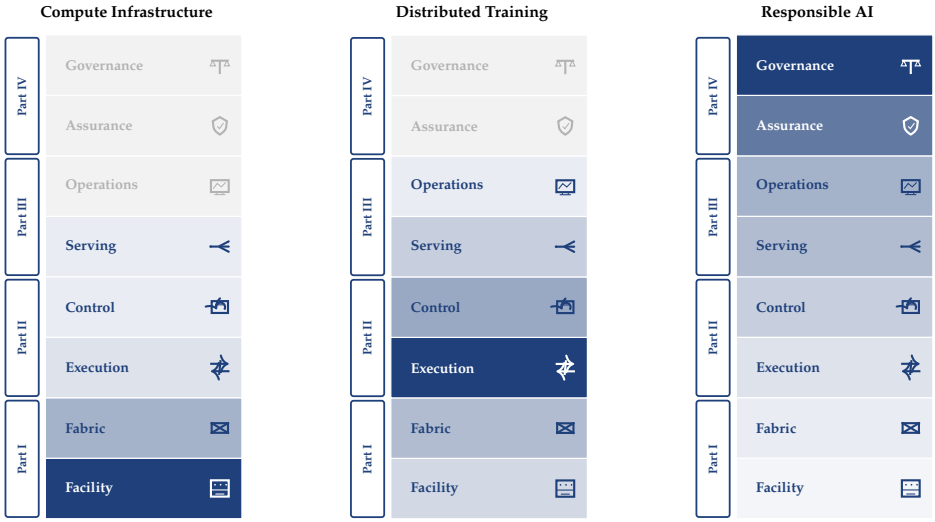

In Compute Infrastructure, the bottom row dominates because the chapter is about silicon, power, cooling, and physical deployment. In Distributed Training, **Distributed Execution** is brightest while **Reliability & Control** remains visibly relevant because large training runs depend on scheduling and recovery. In Responsible AI, **Governance & Sustainability** is primary while **Assurance** stays strong because accountability depends on measurable trust and risk controls. The same stack tells three different stories, and in each one a decision made in one row ripples through the rows above and below it.

Small visual notes sit in the margins wherever a relationship is fastest to grasp by eye. They are reading aids, not figures to cite later. At fleet scale, this visual grammar makes bandwidth gaps, failure blast radius, checkpoint ratios, and edge-energy trade-offs visible without interrupting the argument. When a miniature uses a log scale, it says so inside the image.

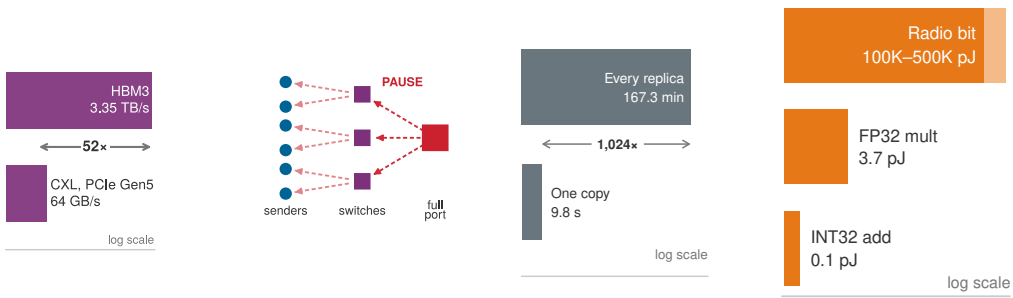

Alongside these visual miniatures, the margins carry unnumbered **marginal navigation pointers** that trace architectural dependencies across the fleet without interrupting the reading flow of the prose:

- ↩ **Backward pointer:** Directs the reader back to the earlier chapter where a physical limit, hardware contract, or analytical model was established.
- ↪ **Forward pointer:** Directs the reader ahead to the later chapter, or to the appendix, where the current idea is developed, measured, or enforced.
- ⇄ **Contrast:** Marks a place where two chapters treat the same problem differently, so the comparison is visible at a glance.

A pointer only orients. When a sentence depends on another chapter, the reference stays in the text.

The appendices are the last reading aid. Beyond the single-machine recap in Appendix A, they collect the diagnostic tables, the communication, reliability, and serving models, and the reference numbers the chapters use.

Conventions

This book uses a small set of typographic conventions to mark what each piece of prose is doing.

Bold marks first definitions

A term in **bold** within body prose is its formal first introduction in the book. Each term receives this treatment exactly once. Subsequent appearances are unbolded; by then the term is part of the working vocabulary the chapter carries forward.

> The **Fleet Law** decomposes the time of one distributed step into local computation, exposed communication, and synchronization.

The bold is a contract: register this term because it will recur. A matching index entry in the back of the book points to the defining occurrence (under the `definition` subentry), providing a way to return to it. Key technical terms are also compiled in the glossary.

Bold also appears in a second, structural role: as a functional label at the start of an item inside a callout box or a summary bullet, for example **Recommendation**, **Trade-off**, or **Result**. In that context, the bold is scaffolding rather than a first definition.

Italics mark specific distinctions

Italics are not used for general emphasis. Every italic span in this book performs one of five specific jobs:

- **Direct opposition**: *training* vs. *inference*, *allocated* vs. *productive* accelerator time.
- **Impossibility, stated as physical, mathematical, or logical law**: "a synchronous step *cannot* finish before its slowest worker does."
- **Word used as word**: "the term *gradient* refers to a vector of partial derivatives."
- **Titles of works**: *Computer Architecture: A Quantitative Approach*.
- **Thesis punchline**: A single italicized sentence that captures the controlling insight of a section, for example, *At scale, failure is the normal operating condition*.

An italicized word is performing one of these jobs. A word that is not italicized carries no implicit emphasis.

Callouts have visual identities that match their job

Callout boxes are not decoration. Each type signals a specific kind of content, so the visual identity tells the reader what the box contains before they read it.

Callout	Contains
Definition	Formal definition of a key term, with its significance, distinction from the nearest neighbor concept, and a common pitfall
Example	Historical case or concrete scenario, structured as context, insight, and systems lesson
Systems Perspective	Analytical observation or mathematical nuance that enriches the surrounding argument
Napkin Math	Worked calculation walked through step by step, with values computed inline from the book's constants; intermediate tables and equations are the math itself
Checkpoint	Self-check questions to verify understanding of the preceding section
War Story	Documented engineering incident with a sourced failure mode, structured as context, failure, consequence, and lesson
Lighthouse	Recurring workloads characterized by a profile (parameters, FLOPs, dominant constraint), used as anchored case studies across multiple chapters
Principle	Canonical book principle or invariant referenced from later chapters

Callout	Contains
Theorem	Formal, equation-backed result or bound, stated with its conditions and applied in later quantitative analysis
Key Takeaways	Four to seven most important results from a chapter, in summary form
Learning Objectives	Skills and concepts that readers should be able to explain or apply after completing the chapter
What's Next	Bridge from the chapter just ended to the chapter beginning, naming the open question the next chapter addresses

When a callout contains a table, figure, or equation, that artifact is part of the callout's pedagogical unit, not decoration. A lighthouse's profile *is* its characterization. A napkin math box's intermediate tables *are* the worked math. Reading these as separated artifacts loses the teaching arc. Standalone reference tables, by contrast, sit between paragraphs as numbered artifacts the reader returns to through cross-references. Both forms are valid; the visual marker (callout vs. standalone) tells the reader which one they are looking at.

Companion Resources

Whichever path readers choose, this book is part of a broader learning ecosystem. The complete text, interactive notebooks, lecture slides, exercises, hands-on labs, and supplementary materials for every chapter are available online at **mlsysbook.ai**.

This book is built in the open. Its full text, figures, and build system are public, and contributions from readers worldwide have sharpened every chapter. Visit **mlsysbook.ai/git** to report an issue, suggest an improvement, or contribute directly.

A Living Textbook

This copy is a snapshot of a book that continues to change. The field keeps moving, and so does my understanding of it, so I keep revising the online edition at **mlsysbook.ai**. The principles in this book are chosen to last. Examples, numbers, and explanations may have improved since this version was published, and the online release notes record what changed.

Using This Book in a Course

This book grew out of CS249r at Harvard University and the TinyML edX professional certificate program. That teaching reinforced an insight that shapes the curriculum here: bandwidth, energy, and failure set the limits on a microcontroller and on a data center alike, and what changes with scale is which limit binds. Because the physics is continuous while the engineering changes at each scale boundary, the book is structured to support a one-semester advanced course covering the full distributed ML systems stack, while allowing instructors to select individual parts for shorter modules:

- **Parts I–II** (The Fleet and Distributed ML) suit an advanced half-semester on distributed training infrastructure.
- **Parts III–IV** (The Fleet in Production and The Responsible Fleet) suit an applied half-semester on production ML at scale. Such a module should first review chapter 5 and chapter 8, which fleet efficiency and serving build on (figure P.1).

For survey courses or executive programs, the four part introductions provide a condensed map of the constraints that govern fleet-scale ML systems. Their principle statements are previews; the chapters supply the assumptions, worked reasoning, and mechanisms needed to apply them.

Slides, assignments, and instructor materials for courses that use this book are available at **mlsysbook.ai**.

Copyright and Licensing

© Vijay Janapa Reddi and MIT Press. This work is distributed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License (CC BY-NC-ND 4.0). The source is available at **mlsysbook.ai/git**.

0

Acknowledgments

Origins

Earlier teaching on TinyML grew out of the realization that the constraints governing milliwatt devices are the same constraints governing data-center accelerators. This book grew out of a different realization, one that came from working on ML Fleet Efficiency as a visiting researcher at Google. Peter Mattson, who had been a collaborator on MLPerf, made that opportunity possible. What I encountered there changed how I think about ML systems: the engineering required to coordinate thousands of accelerators into a single coherent system is a discipline unto itself, and that discipline had no textbook.

That work introduced me to the ML Productivity Goodput (MPG) metric, fleet-level heterogeneity at a scale I had not seen before, and the daily reality of failure as a statistical certainty rather than an exception. The gap between what practitioners knew and what was written down was enormous. The researchers who built these systems published individual results, but the connective tissue between those results, the principles that made them cohere into a discipline, existed only in the heads of the people doing the work. This book is an attempt to document that connective tissue.

Collaborators and Colleagues

I learned most of what I know about fleet-scale systems from the ML Fleet Efficiency team at Google: Arissa Wongpanich, Tayo Oguntebi, Jose Baiocchi Paredes, Yu Emma Wang, Phitchaya Mangpo Phothilimthana, Ritwika Mitra, and Zongwei Zhou. Naveen Kumar led that effort, and the way he thought about fleet-level optimization shaped how I think about it still. Robert Hundt pushed my understanding of performance engineering at scale. Together, Naveen and Robert showed me what fleet-level efficiency looks like from the inside. David Kanter continued the MLPerf collaboration, and our discussions on benchmarking distributed systems sharpened the quantitative methodology that runs through every chapter. Greg Diamos provided invaluable perspective on scaling laws and their origins at Baidu, where some of the earliest work on large-scale distributed training took shape.

The researchers whose papers fill the bibliography of this book deserve particular acknowledgment. The principles taught here were not invented for this textbook; the practitioners and researchers who built the systems, measured the failures, and published what they learned discovered them. This book synthesizes their work into a coherent narrative. The debt is to them.

Support

The Harvard Data Science Initiative, Harvard Extension School, and the National Science Foundation supported this work. The Edge AI Foundation and ICTP supported educational outreach and scholarships, while our industry partners provided the infrastructure and tooling that enabled practical experimentation.

At MIT Press, Susan Hartman believed in this project from the beginning and guided it from an open-source experiment to a published textbook. Her editorial judgment and support made this book real. I am equally grateful to MIT Press for agreeing to keep this book available as open access. A textbook about a discipline shaped by the open-source community should be accessible to everyone in it.

For the complete and most up-to-date list of all GitHub contributors, please visit the online version at mlsysbook.ai. For those interested in contributing, please consult our GitHub repository.

Every GitHub star is a signal that someone cares about this material, is learning from it, and believes the work matters. The online community of learners, with their questions, their curiosity, and their willingness to engage with an unfinished textbook, kept this project moving when the scope felt overwhelming.

Finally, I thank my wife and my children. I wrote this book in the hours that belonged to them: weekends that disappeared into revisions, evenings that stretched past midnight, mornings when I was present but not quite there. They never complained. They asked how the writing was going, brought me coffee, and gave me the space to finish. Whatever merit this book has, their patience and their support made it possible.

0

A Note on AI Assistance

Every chapter of this book was researched, written, and rewritten by me. I also used artificial intelligence tools throughout the process. What follows is more than the disclosure MIT Press requires of authors who use AI. It is an honest statement of how AI assisted in navigating fleet-scale systems engineering, and the rigorous human verification standard applied to the result.

In an era of generative AI, defining what authorship really means has become genuinely difficult. When tools can draft prose, crunch cluster arithmetic, unearth research papers, and generate visual sketches, where does the machine end and the author begin?

In developing this text, AI tools served as an accelerator across computation, literature review, and drafting. On the infrastructure side, this involved extending the MLSysIM infrastructure modeling engine from single-node execution to the fleet level: modeling 3D parallelism (tensor, pipeline, and data parallel strategies), collective communication fabrics, cluster network topologies, and fault-tolerance protocols. AI helped formulate complex cluster-level equations, write executable Python cells, and implement automated pre-commit checks.

AI proved especially valuable in digging into the vast literature of distributed systems, unearthing seminal, classical papers and historical foundations that modern AI literature often overlooks. In the writing itself, I used AI to draft initial chapter outlines, generate preliminary explanations, and synthesize dense research papers into first-pass text, which I then evaluated, restructured, and substantially rewrote to ensure conceptual rigor.

Beyond code and prose, AI fundamentally shifted how I developed the book's visual language. In previous work, sketching architectural diagrams was a slow, manual bottleneck. For this book, I used generative AI to produce initial conceptual drafts and rough schematics of cluster topologies. While these AI-generated sketches were far from publication-ready, they served as an effective conceptual boilerplate, helping clarify and communicate the architectural concepts and communication patterns for the ongoing illustration and refinement process.

Yet mechanical computation, literature synthesis, and AI-assisted drafting are only the starting point. The most critical and time-intensive phase of writing this book was human validation. To ensure that the analytical models reflect the reality of warehouse-scale computing rather than idealized abstractions, I conducted extensive one-on-one review meetings with domain experts: distributed systems researchers, ML fleet operators, and hardware architects. In these sessions, we walked through chapters, stress-tested assumptions, scrutinized performance equations, and debated design trade-offs. This human feedback loop is the longest and most demanding part of the authoring process, but it is indispensable. Without that expert review, a systems book risks teaching clean abstractions that fail on real silicon.

Writing a textbook is an act of synthesis, not a claim of personal invention. The fundamental breakthroughs in distributed systems and hardware architecture belong to the researchers and engineers cited throughout these chapters. My name appears on the cover not out of ego, but because authorship is the assumption of human responsibility. It is the demanding process of vetting these ideas against real-world hardware, and it means providing a single human being whom readers, students, and engineers can hold accountable, someone to reach out to for clarifications, debates, and corrections.

Machine learning systems cannot be certified the way deterministic logic can. That verification gap applies to AI-assisted authoring with equal force. The standard held here is the standard we must demand of every fleet we build: **the tool can propose, but only the human engineer can ratify and take responsibility.**

0

Notation

To a machine learning researcher, B denotes batch size, while to a network engineer, B denotes bandwidth. In distributed systems, N represents the count of worker accelerators, whereas in statistics, N denotes sample count. In high-performance computing, P is commonly used for processor count, but in deep learning, P is reserved for model parameter count. Similarly, in this book's communication models, α denotes per-message network latency and β denotes link bandwidth, whereas in statistical learning, α represents a power-law scaling exponent and β represents optimizer momentum.

These collisions arise because machine learning at scale combines three fields that developed their mathematical conventions independently: machine learning, distributed systems, and high-performance data center networking. Training and serving modern foundation models requires orchestrating workloads across thousands of accelerators connected by high-speed networks, which places symbols from all three fields into the same equation.

This book establishes a unified mathematical notation for distributed machine learning systems. The conventions defined below eliminate ambiguity across scaling laws, collective communication algorithms, parallel execution schedules, and cluster reliability models.

The Governing Equations of Distributed Performance

A distributed training step or serving request spends time both inside each accelerator and across the network that connects them, so performance modeling follows a two-tier hierarchy: the single-device execution foundation and its fleet-wide composition. The subsections that follow state each governing equation and define the symbols that appear in its terms.

The single-device foundation (the iron law of ML systems)

Every distributed training step or inference pass builds on local execution on individual accelerators. For serialized execution phases on a single accelerator, the **iron law of ML systems** decomposes execution time into memory access, arithmetic computation, and latency overhead:

$$T = \underbrace{\frac{D_{\text{vol}}}{\text{BW}}}_{\text{Memory Time}} + \underbrace{\frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}}_{\text{Compute Time}} + \underbrace{L_{\text{lat}}}_{\text{Latency Overhead}}$$

Each symbol was chosen deliberately to preserve machine learning conventions while maintaining architectural precision. The table below records the definition, physical unit, and collision-avoidance rationale for each term:

Symbol	Definition	Unit	Notes
T	**Execution Time**	seconds	Total serialized execution time in the iron-law performance decomposition.
BW	**Bandwidth**	bytes/s	**Avoids collision with B (Batch Size)**. Systems literature often uses B for bandwidth, while ML literature often uses B for batch size. The notation preserves the ML convention.
D_{vol}	**Data Volume**	bytes	**Avoids collision with D (Dataset Size)**. In scaling laws, D means training tokens. Here we need bytes moved through memory. The subscript disambiguates.

Symbol	Definition	Unit	Notes
η_{hw}	**Efficiency**	—	Hardware utilization ($0 \leq \eta_{\text{hw}} \leq 1$). Avoids collision with learning rate (η).
L_{lat}	**Latency**	seconds	**Avoids collision with L (Loss)**. Fixed overhead time (kernel launch, network RTT). The subscript distinguishes from the loss function.
O	**Operations**	FLOPs	Total floating-point operations. Clean in equations (vs. "Ops").
R_{peak}	**Peak Rate**	FLOP/s	**Avoids collision with P (Parameters)**. Roofline presentations often use P for performance, while ML literature commonly uses P for parameter count. The notation preserves the ML convention.
Θ	**Asymptotic Tight Bound**	Operator	Asymptotic tight bound: grows at the same rate as its argument, up to constant factors.
$D_{\text{vol}}^{(\text{batch})}$	**Batch Data Volume**	bytes	Bytes of input data one batch carries to the accelerator.
f	**Clock Frequency**	hertz	Rate at which a processor's clock ticks, setting how many switching cycles occur per second.
Δ	**Finite Change**	Operator	Finite change in the quantity it prefixes, between two states, configurations or points in time.
µ	**Micro Unit Prefix**	Operator	SI unit prefix micro, one millionth, written before a unit such as seconds.
b	**Numeric Bit Width**	Integer	Number of bits used to represent one numeric value, such as an arithmetic operand or a quantized weight or activation.
$T_{\text{pipelined}}$	**Pipelined Execution Time**	seconds	Execution time of a workload whose data movement and computation overlap, set by the slower of the two plus fixed latency.
$T_{\text{sequential}}$	**Sequential Execution Time**	seconds	Execution time when data movement, computation and fixed overhead run one after another with no overlap, the sum of the iron-law terms.
T_{IO}	**Step I/O Time**	seconds	Time the input pipeline needs to read, decompress, decode and stage the data for one training step.
$\text{BW}_{\text{storage}}$	**Storage Bandwidth**	bytes/s	Sustained rate at which the storage system delivers bytes to the input pipeline.
T_{step}	**Training Step Time**	seconds	Wall-clock time of one training step, bounded below by the slower of compute and input I/O.
k	**Unit Alias Multiplier**	Dimensionless	Number of base units in one unit alias, the conversion multiplier of a unit-conventions table row.

Why these choices matter

Without careful notation, statements about distributed systems become ambiguous. Consider a typical unsubscripted claim:

> *"Reducing D improves performance."*

In distributed machine learning, this could mean reducing the dataset size (fewer training samples), shrinking the data-parallel degree N_{DP}, or reducing the data volume D_{vol} moved across the memory bus through quantization. With our notation, we distinguish them precisely:

> *"Reducing D_{vol} through FP8 quantization cuts memory traffic across the interconnect while D (training tokens) and N_{DP} remain unchanged."*

Subscripted variants

The root symbol names the fundamental physical quantity, while an upright subscript specifies the precise domain or measurement point. Three recurring families illustrate this convention across this book:

- $\text{BW}_{\text{memory}}$, $\text{BW}_{\text{network}}$, $\text{BW}_{\text{interconnect}}$: bandwidth across distinct tiers of the distributed memory and fabric hierarchy.
- D_{vol}, R_{peak}, L_{lat}, η_{hw}: single-node iron-law terms, held clear of shared ML symbols.
- E_{move}, E_{compute}, E_{total}: energy budget components across distributed compute nodes.

The subscript therefore carries part of the claim. Memory and network bandwidth differ by orders of magnitude, so a figure is reproducible only when the subscript says where it was measured.

The Fleet Law (distributed step time)

When scaling a model across a fleet of $N = N_{\text{GPU}}$ accelerators, single-device compute time is distributed across workers, while communication and synchronization overheads are introduced. Under partitioned execution of fixed work, the governing equation is the **Fleet Law** (the distributed step-time model, equation 1.3):

$$T_{\text{step}}(N) = \frac{T_{\text{compute}}}{N} + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}}$$

Symbol	Definition	Unit	Notes
N, N_{GPU}	**Number of Devices**	Integer	Accelerator count in a distributed job, the core's N_{GPU} ($N = N_{\text{GPU}}$); also the rank count of a collective. *(Not parameters, use P.)*
G_{node}, N_{nodes}	**Accelerators per Node, Node Count**	Integer	Hierarchical collectives split the fleet as $N = G_{\text{node}} N_{\text{nodes}}$. *(Bare G is the LogGP gap per byte.)*
$T_{\text{step}}(N)$	**Distributed Step Time**	seconds	Wall-clock time for one training step at scale N.
T_{compute}	**Single-Device Compute Time**	seconds	The core's compute term for one training step (forward + backward pass) on one device.
$T_{\text{comm}}(N)$	**Communication Time**	seconds	Time for collective operations (AllReduce, AllGather); depends on payload, topology, and N.
T_{overlap}	**Overlapped Time**	seconds	Communication hidden behind computation. Reduces effective overhead.
$T_{\text{sync}}(N)$	**Synchronization Time**	seconds	Total nonoverlapped synchronization or coordination cost per step at scale N.
ρ	**Communication-Computation Ratio**	Dimensionless	$\rho = T_{\text{comm}}(N)/(T_{\text{compute}}/N)$. Values near or above 1 indicate communication-bound scaling.
CI	**Communication Intensity**	bytes/FLOP	Network bytes moved per local FLOP, CI $= M/O$; the fleet analogue of arithmetic intensity I.
C_2	**Communication Axis**	Set	Communication axis of the C^3 taxonomy, data movement across the network fabric, viewed as the set of techniques it covers so that axis intersections name zones of the taxonomy.
C_1	**Compute Axis**	Set	Compute axis of the C^3 taxonomy, local arithmetic and memory within a node, viewed as the set of techniques it covers so that axis intersections name zones of the taxonomy.
C_3	**Coordination Axis**	Set	Coordination axis of the C^3 taxonomy, synchronization, failure recovery and scheduling across the fleet, viewed as the set of techniques it covers so that axis intersections name zones of the taxonomy.
C	**Operational Dimension**	Record	One of the three operational dimensions of the C^3 taxonomy, indexed 1 to 3 as compute, communication and coordination.

Scaling laws of foundation models

Scaling laws anchor the book's account of why machine learning moved onto fleets. The single-resource power law and the dense-transformer compute rule are:

$$\mathcal{L}(X) = A_{\mathcal{L}} X^{-\alpha_{\text{scale}}} + C_{\mathcal{L}}, \qquad O_{\text{train}} \approx 6PD$$

Symbol	Definition	Unit	Notes
X	**Scaled Resource**	varies	Training compute, parameter count P, or training tokens D, depending on the fitted law.
α_{scale}	**Scaling Exponent**	Dimensionless	Fitted power-law exponent; valid only within the regime in which it was measured. *(Not network latency α.)*
$A_{\mathcal{L}}$	**Scale Coefficient**	loss units	Task-dependent fitted constant.
$C_{\mathcal{L}}$	**Irreducible Loss**	loss units	Loss floor that added resources do not remove within the fitted regime.
O_{train}	**Training Compute**	FLOPs	About $6PD$ for a dense transformer: $2P$ per token forward and $4P$ backward.
B_{base}	**Base Batch Size**	Integer	Reference batch size at which the base learning rate was tuned.
η_{base}	**Base Learning Rate**	Dimensionless	Learning rate tuned at the reference batch size before linear scaling.
N_{epochs}	**Epoch Count**	Integer	Number of full passes a training run makes over its dataset.
B_{new}	**Scaled Batch Size**	Integer	Batch size a training recipe is scaled to under the linear scaling rule.
η_{new}	**Scaled Learning Rate**	Dimensionless	Learning rate used at a new batch size under the linear scaling rule.
$\mathcal{L}$	**Scaling Law Loss**	Function	Loss a trained model reaches as a function of the resource scaled, such as parameters, training tokens or compute, in an empirical power-law scaling law.

The distributed energy corollary

At data center scale, energy and power delivery are primary system constraints. For N devices that each move the same D_{vol} bytes and execute the same O FLOPs over a workload interval, effective per-byte data-movement costs and per-operation compute costs give the following fleet energy model. Here the work and traffic terms are per-device quantities, so multiplying by N sums the device contributions:

$$E_{\text{total}} \approx N \times \left(D_{\text{vol}} \times E_{\text{move}} + O \times E_{\text{compute}} \right)$$

Symbol	Definition	Unit	Notes
E_{move}	**Energy per Byte Moved**	joules/byte	Effective movement cost; total energy also depends on bytes moved and operations executed.
E_{compute}	**Energy per Operation**	joules/FLOP	Energy cost of a single arithmetic operation.
E_{total}	**Total Energy**	joules	Total energy consumed by an ML workload, decomposed into data-movement and compute terms.
i	**Accelerator Supply Current**	amperes	Electrical current an accelerator draws from its power supply.
P_{active}	**Active Inference Power**	watts	Average power an accelerator draws while actively running inference.
P_{avg}	**Average Power**	watts	Power drawn averaged over a stated period.
E_{transmit}	**Cloud Offload Energy**	joules	Energy a device spends transmitting an input to the cloud for remote inference.
P_{thermal}	**Coolant Heat Load**	watts	Heat per unit time that a liquid cooling loop must carry away from a rack, the load that sets its coolant mass flow.
$\dot{m}$	**Coolant Mass Flow Rate**	kg/s	Mass of coolant, liquid or air, passing through a cooling loop per second.
T	**Coolant Temperature**	kelvin	Temperature of the liquid or air coolant in a cooling loop, whose finite change is its rise from inlet to outlet.
P_{device}	**Device Power**	watts	Electrical power one device draws.

Symbol	Definition	Unit	Notes
P_{dynamic}	**Dynamic Power**	watts	Power a chip draws switching its transistors, proportional to switched capacitance, supply voltage squared and clock frequency.
R	**Electrical Resistance**	ohms	Electrical resistance of a circuit or winding.
$\dot{Q}_{\text{remove,max}}$	**Facility Heat Removal Capacity**	watts	Largest rate at which an installed facility and cooling configuration can remove heat.
P_{facility}	**Facility Power**	watts	Total electrical power a data center facility draws from the grid, its IT load plus cooling, power conversion, lighting and other overhead.
$\rho_{\text{energy,general}}$	**General-Purpose Energy Productivity**	FLOP/J	Useful floating-point operations general-purpose silicon delivers per joule on a given workload.
α_{duty}	**Inference Duty Cycle**	Dimensionless	Fraction of time an accelerator actively runs inference rather than waiting in standby.
$\dot{Q}_{\text{IT}}$	**IT Heat Load**	watts	Rate at which a facility's information technology equipment (servers, accelerators, network gear) produces heat.
P_{IT}	**IT Power**	watts	Electrical power delivered to a data center's IT equipment: servers, accelerators, storage and networking.
I_{leak}	**Leakage Current**	amperes	Current that leaks through a chip's transistors when they are not switching.
P_{leakage}	**Leakage Power**	watts	Static power a chip dissipates through leakage current regardless of switching.
E_{local}	**Local Inference Energy**	joules	Energy a device spends running one inference on its own hardware.
E_{network}	**Network Energy**	joules	Energy the network fabric, its switches, links and adapters, consumes over a stated interval.
$L_{\text{parasitic}}$	**Parasitic Loop Inductance**	henries	Unintended inductance of the current loop that feeds an accelerator's core rail, which turns current slew into voltage droop.
$\rho_{\text{energy,specific}}$	**Specialized Energy Productivity**	FLOP/J	Useful floating-point operations specialized, domain-specific silicon delivers per joule on a workload suited to its design.
c_p	**Specific Heat Capacity**	J/(kg K)	Heat needed to raise one kilogram of a coolant by one kelvin at constant pressure.
P_{idle}	**Standby Power**	watts	Power a device draws in standby, when it is not running inference.
V	**Supply Voltage**	volts	Supply voltage that powers a digital circuit, the quantity dynamic voltage and frequency scaling lowers.
C	**Switched Capacitance**	farads	Capacitance of the wires and gates a circuit charges when it switches.
α	**Switching Activity Factor**	Dimensionless	Average fraction of a chip's gate capacitance that switches each clock cycle.
P_{total}	**Total Chip Power**	watts	Total power a chip dissipates, the sum of dynamic switching power and static leakage power.

Network Communication Models and Collectives

Collective communications (AllReduce, AllGather, ReduceScatter, AllToAll) move payloads across the network fabric, following ring or tree schedules over topologies such as the fat tree. The models below price those transfers by separating fixed per-message latency from bandwidth-proportional transmission time.

The α-β communication model

In the standard first-order model, network communication time decomposes into fixed message startup latency and bandwidth-dependent transfer. For a single point-to-point message, transfer time is the sum of those two terms:

$$T(n) = \alpha + \frac{n}{\beta}$$

The model keeps its literature's letters as aliases of core concepts: α is the per-message L_{lat} and β is a link's BW, written in α-β form. Inside α-β formulas, write α and β; elsewhere, write L_{lat} and a subscripted BW.

Symbol	Definition	Unit	Notes
$T(n)$	**Point-to-Point Transfer Time**	seconds	Wall-clock time to transmit one message of size n across one link.
n	**Point-to-Point Message Size**	bytes	Size of one message in $T(n)$.
α	**Per-Message Latency**	seconds	The core's L_{lat} for one message, in α-β form: the fixed startup cost. *(Not significance level or Rényi order, which carry subscripts.)*
β	**Link Bandwidth**	bytes/s	The core's BW for one link (BW_{link}), in α-β form: effective throughput per link.
n^*	**Crossover Point**	bytes	$n^* = \alpha \cdot \beta$. Below n^*: latency-bound. Above: bandwidth-bound.
M	**Collective Payload Size**	bytes	Per-rank gradient, model-state, activation, or compressed payload that a collective moves. Not a memory footprint (those are M_{total}, M_{weights}, and so on) and not the core's GEMM dimension M, which appears only in GEMM shape statements.
M_{coll}	**Collective Tensor Size**	bytes	Full tensor acted on by Broadcast, Reduce, or AllReduce, reconstructed by AllGather, or partitioned by ReduceScatter. An equally sharded AllGather starts with M_{coll}/N bytes per rank; AllToAll uses the per-rank payload M.
T_{ring}, T_{tree}	**Collective Completion Time**	seconds	AllReduce time under the ring or tree algorithm in α-β form, for example $T_{\text{ring}} = 2(N-1)\alpha + 2\frac{N-1}{N}\frac{M}{\beta}$. $T_{\text{allreduce}}$ names an AllReduce whose algorithm the text leaves open.
$M_{\text{crossover}}$	**Ring-Tree Crossover Payload**	bytes	Payload at which $T_{\text{ring}} = T_{\text{tree}}$; approximately $N\alpha\beta/\log_2 N$ for large N.
$T_{\text{allreduce}}$	**AllReduce Completion Time**	seconds	Time to complete one AllReduce collective over a group of devices, whatever algorithm implements it.
T_{alltoall}	**AllToAll Completion Time**	seconds	Time to complete one AllToAll collective, in which every device exchanges a distinct share of its payload with every other device.
$\text{BW}_{\text{bisect}}$	**Bisection Bandwidth**	bytes/s	Total bandwidth across the narrowest cut that divides a network's nodes into two equal halves.
i	**Local Rank Index**	Integer	Position of an accelerator within its node, its local rank, which in a rail-optimized fabric also names the rail it attaches to.
$\text{BW}_{\text{network}}$	**Network Bandwidth**	bytes/s	Sustained rate at which a network path delivers bytes between machines.
M	**Network Traffic Volume**	bytes	Bytes one worker moves across the network over a stated interval, such as one training step.
m	**Rank Count Binary Exponent**	Dimensionless	Integer exponent for which the rank count equals two raised to it, the condition algorithms such as butterfly AllReduce require.
i	**Rank Index**	Integer	Position of one rank, a device or worker, within a collective communication group, counted from zero.
BW_{link}	**Single Link Bandwidth**	bytes/s	Bandwidth of one network link in one direction, written in the core bandwidth notation outside alpha-beta formulas.
k	**Switch Radix**	Integer	Number of ports on one network switch, its radix.
N	**Topology Endpoint Count**	Integer	Number of endpoints, accelerators or hosts, that one network topology connects, the size variable of its scaling relations.

The LogGP extension

LogGP refines the α-β model by separating network transit latency, processor initiation overhead, and long-message injection gaps. The book uses descriptive subscripts for the highest-collision terms; the bare letters g and G keep their LogGP meanings inside this model only.

Symbol	Definition	Unit	Notes
$T_{\text{long}}(n)$	**LogGP Long-Message Time**	seconds	Transfer time for an n-byte long message. Reuses n from the α-β model.
L_{LogGP}	**LogP/LogGP Network Latency**	seconds	Network transit time of one message, the part of the core's L_{lat} that excludes processor overhead, so $\alpha \approx o_{\text{send}} + L_{\text{LogGP}} + o_{\text{recv}}$. The subscript avoids bare L collisions.
o_{send}	**Sender Overhead**	seconds	Processor/NIC overhead to initiate a message.
o_{recv}	**Receiver Overhead**	seconds	Processor/NIC overhead to complete a message. Symmetric LogP writes $o = o_{\text{send}} = o_{\text{recv}}$, so one message costs $2o$.
g	**Message Injection Gap**	seconds	Minimum interval between consecutive message injections. *(Gradients are bold,* $\mathbf{g}$*.)*
G	**Long-Message Gap per Byte**	seconds/byte	Per-byte long-message gap; distinct from β link bandwidth. *(Accelerators per node are* G_{node}*.)*
o	**Symmetric Message Overhead**	seconds	Processor time spent initiating or completing one message in the symmetric LogP model, where sender and receiver overheads are equal, during which the processor cannot compute.

Compression cost and savings

Gradient and activation compression trade local computation for reduced network traffic. The resulting models compare codec overhead against the communication time saved by shrinking the payload:

Symbol	Definition	Unit	Notes
T_{transfer}	**Transfer Time for Payload**	seconds	Communication time for an x-byte payload: $T(x)$ for one link, or the relevant collective time.
$T_{\text{compressed}}$	**Compressed Communication Time**	seconds	Total time after encode, transfer of M/r_{comp}, and decode.
T_{encode}	**Codec Encode Time**	seconds	Local compression overhead before transfer.
T_{decode}	**Codec Decode Time**	seconds	Local decompression overhead after transfer.
b_{block}	**Quantization Block Size**	Integer	Number of tensor elements in one block of block-wise quantization, all sharing one scale.
s	**Quantization Scale**	Varies	Real-valued step between adjacent integer levels of a uniform quantizer.
K_{top}	**Retained Gradient Element Count**	Integer	Number of largest-magnitude gradient elements top-k sparsification keeps and transmits.

Parallelism Dimensions and Scaling Efficiency

Scaling large foundation models requires combining multiple parallelism strategies across device dimensions. The notation below tracks how devices are partitioned across those dimensions and how efficiently the combined schedule converts added hardware into speedup.

Parallelism dimensions

A single parallelism strategy rarely scales across an entire cluster without hitting memory or bandwidth limits. To stay within both bounds, 3D parallelism partitions the training workload across three orthogonal dimensions:

$$N = N_{\text{DP}} \times N_{\text{PP}} \times N_{\text{TP}}$$

Symbol	Definition	Unit	Notes
N_{DP}	**Data-Parallel Degree**	Integer	Number of model replicas. *(Bare d stays the core's hidden dimension.)*
N_{PP}	**Pipeline-Parallel Degree**	Integer	Number of physical pipeline stages in model-depth partitioning.
N_{TP}	**Tensor-Parallel Degree**	Integer	Degree of intra-layer partitioning (model width). *(Bare t stays time or a step index.)*
m	**Microbatch Count**	Integer	Number of microbatches used to fill and drain the pipeline.
b_{pipe}	**Pipeline Bubble Fraction**	Dimensionless	Fraction of pipeline GPU-time spent idle, $b_{\text{pipe}} = (N_{\text{PP}} - 1)/(N_{\text{PP}} - 1 + m)$; avoids bare b.
V	**Virtual Pipeline Stages per Device**	Integer	Number of interleaved virtual chunks assigned to each physical pipeline stage. *(The attention value matrix is bold, $\mathbf{V}$.)*
τ_{stale}	**Gradient Staleness**	steps	Number of parameter updates between the parameters a gradient was computed on and the parameters it updates; $\tau_{\text{stale}} = 0$ is synchronous training.
k	**Accumulation Step Count**	Integer	Number of micro-batches whose gradients are accumulated before one optimizer update.
k	**Client Index**	Integer	Position of one client among the clients of a federated round or secure aggregation.
C_{cap}	**Expert Capacity Factor**	Dimensionless	Multiplier on an expert's uniform share of token assignments that sets how many assignments it may accept per step.
N_{experts}	**Expert Count**	Integer	Number of expert subnetworks in one mixture-of-experts layer.
R	**Federated Round Count**	Integer	Number of communication rounds a federated training run performs, or needs to reach a quality target.
B_{global}	**Global Batch Size**	Integer	Number of examples one synchronous optimizer update averages over across all data-parallel workers.
$\mathbf{g}$	**Local Gradient**	Vector	Gradient vector one data-parallel worker computes from its local batch in one step, its contribution to the gradient collective.
B_{micro}	**Microbatch Size**	Integer	Number of examples in one micro-batch, the unit run through a single forward and backward pass.
N_{clients}	**Participating Client Count**	Integer	Number of clients whose updates one federated round aggregates.
k_{experts}	**Routed Expert Count**	Integer	Number of experts a mixture-of-experts router sends each token to.
p_{gate}	**Router Gate Probability**	Dimensionless	Probability a mixture-of-experts router assigns to the expert it selects for a token.

Scaling efficiency

Adding accelerators to a job increases communication and synchronization overhead alongside raw compute capacity. Scaling efficiency measures the resulting realized speedup relative to ideal linear scaling:

$$\eta_{\text{scaling}} = \frac{T_1}{N \times T_N}$$

Symbol	Definition	Unit	Notes
η_{scaling}	**Scaling Efficiency**	Dimensionless	Ratio of measured speedup to ideal linear speedup. 1.0 is the ideal linear baseline, not a hard upper bound.
T_1	**Single-Device Time**	seconds	Baseline wall-clock time for the same productive work on one device.
T_N	**N-Device Time**	seconds	Wall-clock time for the same productive work on N devices; excludes operational downtime counted by η_{runtime}.
r	**Serial Fraction**	Dimensionless	Share of a fixed workload's single-device execution time that does not shrink when the work is spread across devices, the serial term of Amdahl's law.

Cluster Reliability, Fault Tolerance, and Checkpointing

At cluster scale, hardware component failures are standard operating conditions rather than rare anomalies. The symbols in this section quantify failure rates, checkpoint intervals, recovery overheads, and their net effect on fleet goodput.

Fault tolerance and reliability

When a job requires many components to remain healthy simultaneously, individual failure hazards compound across the fleet. Assuming independent components with exponential failure distributions, system reliability across N components is:

$$R_{\text{system}}(t) = e^{-N\lambda t}$$

The aggregate exponent uses the system-level rate $N\lambda$, where λ is the per-component failure rate under the independent, identical-component assumption. The exponent $N\lambda t$ is dimensionless, so the elapsed time t must use the same time unit as λ (hours, when λ comes from the FIT-derived rates below).

Checkpointing too frequently wastes wall time writing state to storage, whereas checkpointing too rarely loses long stretches of training work when a node fails. The optimal checkpoint interval balances that I/O overhead against recomputation rework:

$$T^*_{\text{ckpt}} = \sqrt{2 \cdot T_{\text{write}} \cdot \text{MTBF}_{\text{system}}}$$

The formula uses the *system-wide* mean time between failures (the failure interval for the whole job), not a *component* lifetime; substituting the per-component value would overestimate T^*_{ckpt} by $\sqrt{N}$. For dimensional consistency, T_{write} and $\text{MTBF}_{\text{system}}$ must be in the same time unit. Industry-standard MTBF is reported in hours; to combine it with seconds-scale write times, multiply the hour value by 3600. The result T^*_{ckpt} then carries the time unit chosen for the inputs.

Symbol	Definition	Unit	Notes
$R_{\text{system}}(t)$	**Reliability Function**	Probability	Probability of no failure across the whole system before time t. One component has $R_{\text{single}}(t) = e^{-\lambda t}$.
λ	**Per-Component Failure Rate**	failures/hour	Per-component failure rate, written λ_{fail} wherever another rate shares the passage; the system-wide rate is $N\lambda$ under independent, identical components. The industry-standard reporting unit is FIT (failures per 10^9 device-hours): $\lambda_{\text{hour}} = \text{FIT} \times 10^{-9}$. *(Other rates carry subscripts, such as the core's arrival rate λ_{arr}.)*
$A_{\text{component}}$	**Component Availability**	Probability	Probability that one component is up at a given instant. A job that needs all N components at once has availability $(A_{\text{component}})^N$ under independent, identical components.
N	**Required Component Count**	Integer	Components (accelerators, hosts, NICs, switches) the job needs up at the same time. Not N, which counts accelerators only.

Symbol	Definition	Unit	Notes
MTBF	**Mean Time Between Failures**	hours	Average time between consecutive failures of a repairable, composed system. $\text{MTBF}_{\text{system}} = 1/(N\lambda) = \text{MTTF}_{\text{component}}/N$.
$\textbf{MTBF}_{\textbf{system}}$	**System MTBF**	hours	Mean time between failures that interrupt the whole job, equal to the per-component MTTF divided by N under independent, identical components. It is the input to the Young-Daly interval. Same time unit as T_{write} when combined with it.
MTTF	**Mean Time To Failure**	hours	Per-component lifetime constant derived from FIT rates, $\text{MTTF}_{\text{component}} = 1/\lambda$. Distinct from system-level MTBF when repair/replacement cycles are modeled.
MTTR	**Mean Time To Repair**	hours	Average recovery time after a failure.
T_{ckpt}	**Checkpoint Interval**	seconds	Useful compute time between consecutive checkpoints.
T^{*}_{ckpt}	**Optimal Checkpoint Interval**	seconds	The T_{ckpt} that minimizes total wasted time; Young-Daly formula.
T_{write}	**Checkpoint Write Time**	seconds	Time to persist model state to storage.
T_{load}	**Checkpoint Load Time**	seconds	Time to read the latest checkpoint state from storage back into accelerator memory during recovery.
$\textbf{MTTF}_{\textbf{component}}$	**Component Mean Time to Failure**	hours	Mean time to failure of one component, the reciprocal of its constant failure rate.
R_{single}	**Component Reliability**	Dimensionless	Probability that one component of a system runs without failure up to a given time, its survival function.
d	**Erasure Data Fragment Count**	Integer	Number of data fragments an erasure code splits an object into, any that many of the stored fragments recovering the original.
m	**Erasure Parity Fragment Count**	Integer	Number of parity fragments an erasure code adds to the data fragments of an object.
T_{detect}	**Failure Detection Time**	seconds	Time from a hardware or process fault until the system classifies it as a failure.
T_{recovery}	**Failure Recovery Time**	seconds	Time from a failure until useful training work resumes: detection, restart, checkpoint load and warmup.
$\textbf{MTTF}_{\textbf{GPU}}$	**GPU Mean Time to Failure**	hours	Mean time to failure of one GPU, its average operating time before it fails.
A_{all}	**Job Availability**	Dimensionless	Probability that every component a job needs is up at the same instant.
T_{restart}	**Job Restart Time**	seconds	Time for the scheduler to allocate resources and launch replacement processes after a failure.
T_{replay}	**Lost Work Replay Time**	seconds	Time to recompute the training steps lost since the last checkpoint after a restart.
k	**Redundant Replica Count**	Integer	Number of independent replicas of a service in a redundancy availability calculation, any one of which can serve the load.
A_{system}	**Redundant Service Availability**	Dimensionless	Probability that at least one of a service's independent replicas is up at a randomly chosen instant.
A_{single}	**Replica Availability**	Dimensionless	Probability that one serving replica is up at a randomly chosen instant.
n_{psu}	**Series Power Supply Count**	Integer	Number of power supply units in one node whose failure alone stops the node.
A	**Steady-State Availability**	Dimensionless	Probability that a component, replica or repairable system is up at a randomly chosen instant, its long-run fraction of up time.

Fleet capacity factors and goodput

Fleet capacity metrics compound single-device arithmetic efficiency, scaling efficiency, and operational losses. ML productivity goodput multiplies the scheduling, runtime, and program factors that chapter 9 defines.

Symbol	Definition	Unit	Notes
f_{compute}	**Useful Compute-Time Fraction**	Dimensionless	$(T_{\text{compute}}/N)/T_{\text{step}}(N)$. Measures how much step time is useful local arithmetic after distribution.
f_{overhead}	**Operational Overhead Fraction**	Dimensionless	Fraction of allocated wall time lost outside the productive steps measured by T_N, such as checkpoint pauses, failure recovery, scheduler preemption, or maintenance.
R_{eff}	**Effective Fleet Throughput**	FLOP/s	$R_{\text{eff}} = (NR_{\text{peak}}) \times \text{MFU} \times \eta_{\text{scaling}} \times \eta_{\text{runtime}}$, where MFU is the single-device baseline, η_{scaling} captures all in-step scaling loss, and η_{runtime} captures only out-of-step operational loss excluded from T_N.
O_{useful}	**Useful Model FLOPs**	FLOPs	Model FLOPs that advance training, excluding recomputation and padding; $\text{MFU} = O_{\text{useful}}/(NR_{\text{peak}}T_{\text{step}})$ measures fleet MFU and already includes in-step scaling loss. The throughput factorization uses single-device baseline MFU instead.
MPG, η_{program}, η_{runtime}, $\eta_{\text{scheduling}}$	**ML Productivity Goodput and Its Factors**	Dimensionless	Program, runtime, and scheduling goodput, with $\text{MPG} = \eta_{\text{scheduling}} \times \eta_{\text{runtime}} \times \eta_{\text{program}}$. For a job that holds its devices, R_{eff} equals peak capacity times $\eta_{\text{runtime}} \times \eta_{\text{program}}$, so η_{runtime} is also the operational factor of that product; MPG adds the scheduling factor.
ρ_{energy}	**Fleet Energy Productivity**	FLOP/J	Useful computation per unit energy, O_{useful} divided by the energy drawn inside the stated boundary.
C_{train}, C_{infer}	**Dollar Cost**	dollars	The subscript names the cost component. Carbon quantities use $\text{CO}_{2,\ldots}$, never C.
T_{wall}	**Allocated Wall Time**	seconds	Total wall-clock time for which a training job holds its allocated resources, productive or not.
T_{job}	**Job Duration**	seconds	Wall-clock time one job or run spends executing, from start to finish.
T_{useful}	**Useful Training Time**	seconds	Wall-clock time a training job spends on committed, productive steps that contribute to the final model.

Distributed Memory Hierarchy and Serving at Scale

Managing distributed memory requires partitioning weights, activations, optimizer states, and KV-cache tensors across nodes and tiers. At the same time, serving fleets must bound queueing delay and batching overhead under variable request arrivals.

Distributed memory and serving metrics

Reusable performance and serving quantities follow the same collision-avoidance rule as the iron law. Rates use R or descriptive Greek symbols; request counts use Q_{req} rather than overloading N; queue utilization uses a subscripted ρ so bare ρ remains available for other ratio models.

Symbol	Definition	Unit	Notes
$M_{\text{activations}}$	**Activation Memory**	bytes	Retained activations for the backward pass or serving intermediates.
λ_{arr}	**Arrival Rate**	requests/s	Request arrival rate. The subscript keeps bare λ free for fitted sensitivities.
R_{attain}	**Attainable Compute Rate**	FLOP/s	Roofline bound: $R_{\text{attain}} \leq \min(R_{\text{peak}}, I \times \text{BW})$. Uses R, not T, because this quantity is a rate.
$T_{\text{svc}}(B)$	**Batch Service Time**	seconds	Time to serve a batch of size B.
r_{comp}	**Compression Ratio**	Dimensionless	Uncompressed size divided by compressed size. A compressed payload has $1/r_{\text{comp}}$ of the uncompressed size; the subscript avoids bare C collisions.
$\mu_{\text{eff}}(B)$	**Effective Service Rate**	requests/s	Batched service rate, typically $\mu_{\text{eff}}(B) = B/T_{\text{svc}}(B)$. The unbatched service rate is $\mu = \mu_{\text{eff}}(1)$.
s_{elem}	**Element Storage Size**	bytes/element	Bytes per stored tensor element.
$M_{\text{gradients}}$	**Gradient Memory**	bytes	Memory occupied by stored gradients.

Symbol	Definition	Unit	Notes
MFU	**Model FLOPs Utilization**	Dimensionless	Useful model FLOP/s divided by available peak FLOP/s. Text acronym avoids overloading η.
$M_{\text{optimizer}}$	**Optimizer-State Memory**	bytes	Momentum, variance, master weights, and related optimizer buffers.
R	**Replica Count**	Integer	Serving replicas among which a router chooses (for example, probing two of R and sending to the less loaded). Distinct from N, the devices in one job.
Q_{req}	**Request Concurrency**	requests	Average in-flight requests in a stable serving system ($Q_{\text{req}} = \lambda_{\text{arr}} \cdot T_{\text{lat}}$ via Little's Law). Avoids overloading N, which counts devices, elements, or samples elsewhere.
T_{lat}	**Request Time in System**	seconds	End-to-end queueing/serving latency for Little's Law. Distinct from L_{lat}, the fixed-latency term in the iron law.
I	**Arithmetic Intensity**	FLOP/byte	Workload FLOPs per byte moved. The Roofline Model uses I as the independent variable.
I_{ridge}	**Roofline Ridge Point**	FLOP/byte	$I_{\text{ridge}} = R_{\text{peak}}/\text{BW}$. Prefer this explicit form over starred shorthand so the meaning remains clear in prose.
ρ_{serv}	**Serving Utilization**	Dimensionless	Queue/server utilization. The subscript keeps bare ρ free for other ratio models.
M_{total}	**Total Memory Footprint**	bytes	Sum of explicit memory components; avoids bare M ambiguity; the core's quantity.
M_{weights}	**Weight Memory**	bytes	Memory occupied by model weights ($P \times$ bytes-per-parameter); the core's quantity.
BW_{mem}	**Sustained Memory Bandwidth**	bytes/s	Empirically achievable memory bandwidth under target access patterns.
$R_{\text{consumption}}$	**Accelerator Data Consumption Rate**	bytes/s	Rate at which one accelerator consumes training input data while computing at full speed.
λ_{batch}	**Batch Arrival Rate**	batches/s	Rate at which batches arrive at a server when each batch is treated as one super-request, the request arrival rate divided by the batch size.
k	**Batch Fill Arrival Count**	Integer	Number of arrivals after a batch's triggering request needed to fill the batch, one less than the maximum batch size.
T_{fixed}	**Batch Fixed Overhead Time**	seconds	Part of a batch's service time paid once per batch whatever its size, such as weight loading and kernel launch.
$L_{\text{lat,form}}$	**Batch Formation Delay**	seconds	Time a request waits for its batch to fill before the batch is dispatched.
k	**Batch Index**	Integer	Position of one batch in the sequence of batches a data loader or runtime delivers.
T_{window}	**Batching Window**	seconds	Length of the batching window, the longest a dynamic batcher holds requests before dispatching a batch.
Q_{day}	**Daily Query Volume**	queries/day	Number of queries, requests or transactions a deployed model serves per day.
I_{decode}	**Decode Arithmetic Intensity**	FLOP/byte	Arithmetic intensity of the weight-dominated part of an LLM decode step, operations per byte of weights and activations moved, as a function of batch size.
M_{device}	**Device Memory Capacity**	bytes	Total memory a device provides for a model's weights and activations.
α_{acc}	**Draft Token Acceptance Probability**	Dimensionless	Probability that the target model accepts one drafted token in speculative decoding, taken as independent across draft positions.
$L_{\text{lat,total}}$	**End-to-End Request Latency**	seconds	Time from a request leaving the client to its response returning, summed over every stage on the request path.
B	**Error Budget Consumption**	Dimensionless	Share of one reliability target's error budget a service has used: the fraction of eligible observations that violate the target divided by the violating fraction the target allows.
$L_{\text{lat,infer}}$	**Inference Latency**	seconds	Time the model takes to compute a request's output.
$L_{\text{lat,SLO}}$	**Latency Service-Level Objective**	seconds	Latency target a serving system's service-level objective sets for a request.
t_{req}	**Marginal Request Service Time**	seconds	Additional batch service time that one more request in the batch adds.
B_{max}	**Max Batch Size**	Integer	Largest batch a dynamic batcher forms, the size at which it dispatches without waiting for the window to close.
B_{min}	**Minimum Stable Batch Size**	Integer	Batch size at which a batched server's effective service rate equals the arrival rate, the boundary that stable batch sizes must exceed.

Symbol	Definition	Unit	Notes
$L_{\text{lat,network}}$	**Network Latency**	seconds	Time a request and its response spend crossing the network, including propagation, routing and protocol handling.
N_{prefetch}	**Prefetch Buffer Count**	Integer	Number of buffers holding batches loaded ahead of use.
c	**Queue Server Count**	Integer	Number of identical servers in a multi-server queue with random arrivals and service.
K	**Queue System Capacity**	Integer	Largest number of requests a queueing system can hold at once, waiting plus in service, beyond which arrivals are rejected.
W_q	**Queue Wait Time**	seconds	Mean time a request waits in a queue before its service begins, excluding service time.
$L_{\text{lat,queue}}$	**Queueing Latency**	seconds	Time a request waits in queues before it is served.
n	**Realized Arrival Count**	Integer	One possible value of the number of requests that arrive after the trigger in a batching window, used as a summation variable.
i	**Reliability Target Index**	Integer	Position of one reliability target, such as availability, latency, freshness, or model quality, among the targets an evaluation contract budgets.
t	**Request Response Time**	seconds	Time one request spends in a serving system, the random variable whose distribution gives tail latency.
σ_t	**Response Time Standard Deviation**	seconds	Standard deviation of the time a request spends in a serving system.
μ	**Service Rate**	items/s	Rate at which one server, such as a worker, consumer, or model endpoint, completes items or calls.
K	**Speculative Draft Length**	Integer	Number of tokens a draft model proposes for the target model to verify in one speculative-decoding round.
τ	**Subsequent Arrival Time**	seconds	Time from the request that opened a batching window until a later arrival, indexed by arrival order.
$T_{\text{lat,p99}}$	**Tail Request Latency (99th Percentile)**	seconds	99th-percentile time a request spends in a serving system.
ρ_{target}	**Target Serving Utilization**	Dimensionless	Serving utilization a system is sized not to exceed, chosen below one to keep latency headroom.
B_{util}	**Target Utilization Batch Size**	Integer	Smallest batch size whose serving utilization stays at or below the target utilization, under the linear batch service model.
Q_{arr}	**Triggered Window Arrival Count**	Integer	Number of requests that arrive after the request that opened a batching window and before the window closes.

Serving memory and data storage

At scale, serving is bounded by the memory its key-value cache occupies, and the data behind training is bounded by what it costs to keep. These symbols size both.

Symbol	Definition	Unit	Notes
M_{KV}	**KV-Cache Memory**	bytes	Total key-value cache footprint for autoregressive serving, $M_{\text{KV}} = B \cdot S \cdot m_{\text{token}}$.
m_{token}	**KV-Cache Bytes per Token**	bytes/token	$m_{\text{token}} = 2N_L H_{\text{KV}} d_{\text{head}} s_{\text{elem}}$ (keys and values across all layers).
C_{store}	**Data Storage Cost**	dollars	Cost of keeping a training dataset in storage over its retention period.
N_{entities}	**Entity Count**	Integer	Number of entities, such as users or items, a recommendation table or feature store holds rows for.
T_{months}	**Storage Retention Time**	months	Length of time a dataset is kept in storage.
$D_{\text{vol,store}}$	**Stored Data Volume**	bytes	Bytes a retained dataset occupies in storage.
C_{total}	**Total Data Cost**	dollars	Total cost of a training dataset over its lifecycle: acquisition, labeling, storage and processing.

Deep Learning Foundations in Distributed Contexts

Distributed algorithms partition neural architectures defined by standard deep learning conventions (Goodfellow et al. 2016):

Symbol	Definition	Unit	Notes
d	**Hidden Dimension**	Integer	Size of the hidden state vector. *(Parallelism degrees carry subscripts, such as N_{DP}.)*
d_{head}	**Attention Head Dimension**	Integer	Per-head hidden dimension in attention layers.
d_{in}	**Layer Input Width**	Integer	Number of input features in a layer.
d_{out}	**Layer Output Width**	Integer	Number of output features in a layer.
θ	**Model Parameters**	Vector	Trainable model weight and bias parameter vector θ.
π_θ	**Policy**	Function	Target neural network policy $\pi_\theta(a \mid s)$.
S	**Sequence Length**	Integer	Input context or sequence length S in tokens.
ℓ	**Layer Index**	Integer	Index for a layer. Use instead of bare L when indexing layers, since L collides with loss and latency. A numeric subscript marks a norm instead (ℓ_1, ℓ_2).
$\mathcal{L}$	**Loss Function**	Scalar	The objective function minimized during training.
$N_{\text{heads}}, H_{\text{KV}}$	**Number of Attention Heads**	Integer	Number of attention heads in a multi-head attention layer, and key-value heads H_{KV} (in GQA/MQA).
N_L	**Number of Layers**	Integer	Total number of layers in a network.
$p(x)$	**Distribution**	Probability mass/density	Use lowercase $p(\cdot)$ for generic distributions to avoid collision with P (parameter count).
$p(y \mid x)$	**Conditional Distribution**	Conditional probability/density	Use this form for generic label relationships.
$\Pr(\cdot)$	**Event Probability**	Probability	Probability of an event E, as in $\Pr(\text{batch} = 0)$; use $p(x)$ and $p(y \mid x)$ for distributions.
m	**Adam First Moment**	Varies	Per-parameter exponential moving average of gradients that Adam keeps, its first moment.
v	**Adam Second Moment**	Varies	Per-parameter exponential moving average of squared gradients that Adam keeps, its second moment.
$\mathbf{A}$	**Adapter Down Projection Matrix**	Matrix	First low-rank adapter matrix, which projects a layer's input down to the adapter rank, the LoRA down-projection.
r	**Adapter Rank**	Integer	Rank of a low-rank adapter update, equivalently the bottleneck width of a residual adapter, which sets its trainable parameter count.
$\mathbf{B}$	**Adapter Up Projection Matrix**	Matrix	Second low-rank adapter matrix, which projects the rank-sized intermediate up to the layer's output width, the LoRA up-projection.
$\mathbf{K}$	**Attention Key Matrix**	Matrix	Matrix of key vectors, one row per token, in an attention layer.
$\mathbf{Q}$	**Attention Query Matrix**	Matrix	Matrix of query vectors, one row per token, in an attention layer.
S_{ij}	**Attention Score Entry**	Dimensionless	Attention score of query i against key j, one entry of the score matrix.
$\mathbf{S}$	**Attention Score Matrix**	Matrix	Matrix of attention scores, the scaled dot products of every query with every key before the softmax.
$\mathbf{V}$	**Attention Values Matrix**	Matrix	Matrix of value vectors, one row per token, that attention weights combine into the output.
$\mathbf{b}$	**Bias Vector**	Vector	Vector of learned biases one layer adds to its weighted sums, one per neuron or output channel.
K	**Class Count**	Integer	Number of classes a classifier distinguishes, or vocabulary entries it scores, the length of its softmax output.
j	**Class Index**	Integer	Index of an output class.
$\mathbf{z}$	**Encoder Latent Representation**	Vector	Intermediate representation a shared encoder computes from an input, consumed by the heads that sit on top of it.

Symbol	Definition	Unit	Notes
$\mathcal{N}$	**Gaussian Distribution**	Distribution	Gaussian (normal) distribution with the mean and covariance given as its arguments.
A	**GEMM Left Operand**	Matrix	Left input matrix of a general matrix multiply, whose rows match the output and whose columns match the right operand's rows.
C	**GEMM Output Matrix**	Matrix	Output matrix of a general matrix multiply, with the left operand's row count and the right operand's column count, which the operation updates in place.
B	**GEMM Right Operand**	Matrix	Right input matrix of a general matrix multiply, whose rows match the left operand's columns.
M	**GEMM Row Dimension**	Integer	Number of rows of the left operand and of the output in a matrix multiply.
h	**Hidden Activation**	Vector	Activation vector a hidden layer outputs and passes to the next layer.
h	**Hidden Unit Output**	Dimensionless	Output of one hidden unit of a small network.
I	**Identity Matrix**	Matrix	Square matrix with ones on the diagonal and zeros elsewhere.
k	**Key Vector**	Vector	Key vector of one token, one row of the key matrix.
j	**Key-Value Tile Index**	Integer	Index of one key-value tile in tiled attention.
i	**Label Category Index**	Integer	Position of one label category among a classifier's output classes.
ℓ	**Layer Index**	Integer	Position of one layer within a network, counted from the input.
X	**Layer Input Matrix**	Matrix	Input to a layer in matrix form, one row per example or token, or a single-channel image.
x	**Layer Input Vector**	Vector	Vector of activations one layer receives as input.
Y	**Layer Output Matrix**	Matrix	Output matrix of a layer's matrix multiplication of the input matrix by the weight matrix, one row per example.
y	**Layer Output Vector**	Vector	Vector of activations one layer produces.
g	**Gradient**	Varies	Gradient of the training loss with respect to one model parameter, as an optimizer update uses it.
H	**Loss Hessian**	Matrix	Matrix of second derivatives of the loss with respect to every pair of model parameters.
D	**Multiply-Accumulate Destination**	Matrix	Destination matrix of a matrix multiply-accumulate instruction, which receives the product plus the accumulator input.
i	**Matrix Row Index**	Integer	Position of a row within a matrix.
f	**Model Predictor**	Function	Trained model viewed as a function that maps an input to its prediction.
s	**Model Score**	Dimensionless	Score a classifier assigns to one individual, read as the predicted probability of the positive outcome.
w	**Model Weight**	Dimensionless	Value of one learnable weight of a model, such as the weight on one input of a neuron.
y	**Network Output**	Dimensionless	Scalar output of a neuron or of a small network.
θ^*	**Optimal Model Parameters**	Vector	Model parameters that minimize the training loss for a given architecture.
t	**Optimizer Step Index**	Integer	Position of one optimizer update in the sequence of training steps.
p	**Predicted Class Probability**	Dimensionless	Probability a classifier assigns to one label category for a given input.
$\hat{\mathbf{y}}$	**Predicted Output**	Vector	Model output for one example, such as its vector of class probabilities.

Symbol	Definition	Unit	Notes
q_{im}	**Query Component**	Dimensionless	One component of one token's query vector.
$\mathbf{q}$	**Query Vector**	Vector	Query vector of one token, one row of the query matrix.
x	**Raw Pipeline Input**	Record	Raw input record before feature transformation, as received by the training or serving pipeline.
f_{serve}	**Serving Transform**	Function	Transformation the serving pipeline applies to a raw input to produce model features.
N	**Square GEMM Dimension**	Integer	Side length of the square matrices in a matrix multiply, the shared row, column and inner dimension.
Φ	**Standard Normal Cumulative Distribution**	Function	Cumulative distribution function of the standard normal distribution.
k	**Top Class Count**	Integer	Number of highest-scoring classes a top-k prediction returns or checks for the correct label.
f_{train}	**Training Transform**	Function	Transformation the training pipeline applies to a raw input to produce model features.
$\mathbf{W}$	**Weight Matrix**	Matrix	Matrix of learned weights one layer multiplies its input by, or the kernel tensor of a convolution layer.
n	**Weight Matrix Column Count**	Integer	Number of columns of a layer's weight matrix.
m	**Weight Matrix Row Count**	Integer	Number of rows of a layer's weight matrix.
λ_{reg}	**Weight Penalty Strength**	Dimensionless	Coefficient that scales an absolute-value or squared penalty on weight magnitude added to the training loss, equivalent to weight decay under plain gradient descent.
$\mathbf{0}$	**Zero Vector**	Vector	Vector whose entries are all zero, such as the mean of zero-centered Gaussian noise.

Latin matrices and vectors are set in bold (**W**, **x**); scalars, dimensions, indices, and individual matrix entries stay italic (N, d, i, W_{ij}). The generic parameter symbol θ is the conventional exception and stays italic. Matrix operands in general matrix multiply (GEMM) are bold ($\mathbf{C} = \alpha\mathbf{AB} + \beta\mathbf{C}$), distinct from the scalar batch size B. GEMM dimensions are M, N, and K. Asymptotic complexity is written $\mathcal{O}(\cdot)$, distinct from total operations O.

Privacy, Robustness, Governance, and Distributed Constants

Federated systems, distributed coordination, and cluster operations add the statistical, privacy, and governance quantities collected below. These symbols standardize how the later chapters express differential privacy budgets, adversarial perturbation radii, drift statistics, and operational deadlines.

Privacy, robustness, and statistical testing

Part IV uses the standard symbols of differential privacy and adversarial robustness, and the operations and robustness chapters use those of hypothesis testing. The same Greek letters appear in several of these fields, so every use other than the one named below carries a subscript.

Symbol	Definition	Unit	Notes
(ϵ, δ)	**Differential Privacy Budget**	Dimensionless	ϵ bounds the privacy loss from adding or removing one protected record; δ is the allowed failure probability.
ϵ_{adv}	**Adversarial Perturbation Budget**	input units	Radius of the allowed input perturbation under a stated norm. *(Not the privacy budget ϵ.)*
K_{PGD}	**Attack Steps**	Integer	Inner iterations of a projected gradient descent (PGD) attack; adversarial training costs about $K_{\text{PGD}} + 1$ passes per batch.

Symbol	Definition	Unit	Notes
α_{sig}, β_{stat}	**Type I and Type II Error Rates**	Probability	Significance level and miss rate of a hypothesis test; statistical power is $1-\beta_{\text{stat}}$.
$\mathbb{E}[\cdot]$	**Expectation**	—	Expected value. Blackboard bold keeps it distinct from an event E and the energy terms such as E_{total}.
x_{adv}	**Adversarial Example**	Record	Adversarial example, an input an attacker forms by adding a bounded perturbation to a clean input so the model mispredicts.
δ_{adv}	**Adversarial Perturbation**	Vector	Perturbation an adversary adds to an input to change a model's prediction.
k_{anon}	**Anonymity Set Size**	Integer	Smallest number of records that share every quasi-identifier value in an anonymized release.
n_{feat}	**Attribution Feature Count**	Integer	Number of input features a feature-attribution method assigns contributions to.
γ_{cls}	**Classification Threshold**	Dimensionless	Score threshold above which a classifier outputs the positive class.
$z_{\alpha/2}$	**Confidence Critical Quantile**	Dimensionless	Positive standard normal critical value for a two-sided test at the chosen confidence level.
ϵ	**Conformal Miscoverage Rate**	Dimensionless	Allowed probability that the conformal prediction set misses the true referent.
$\mathcal{C}$	**Conformal Prediction Set**	Set	Calibrated set of referent candidates that cannot be excluded at the miscoverage rate.
A	**Current Bin Proportion**	Dimensionless	Proportion of current production observations falling in one bin of the population stability index.
P_t	**Current Distribution**	Distribution	Data distribution a deployed model sees in service at time t.
n	**Drift Test Sample Size**	Integer	Number of observations in one of the two samples a two-sample drift test compares.
g	**Environment Response Function**	Function	Response of the environment to a deployed model, mapping the current data distribution and the model's decisions to the next data distribution.
ϵ_{fair}	**Fairness Gap Tolerance**	Dimensionless	Largest allowed gap between protected groups in the fairness metric a release gate or monitor enforces.
t	**Feature Record Time**	seconds	Time at which a feature record became available, which a point-in-time join requires to be no later than the observation timestamp.
a	**First Compared Group Label**	Record	First of two group labels compared in a group-fairness condition, one value of the protected group attribute.
$\mathcal{D}_{\text{JS}}$	**Jensen-Shannon Divergence**	Function	Jensen-Shannon divergence between two probability distributions, a symmetric bounded measure of how far they differ.
$\bar{P}$	**Jensen-Shannon Mixture Distribution**	Distribution	Equal-weight average of the current and training distributions, the reference against which Jensen-Shannon divergence measures each.
$\mathcal{D}_{\text{KL}}$	**KL Divergence**	Function	Kullback-Leibler divergence, the expected extra code length from encoding one distribution with a code built for another.
$\mathcal{D}_{\text{KS}}$	**Kolmogorov Smirnov Statistic**	Dimensionless	Two-sample Kolmogorov-Smirnov statistic, the largest absolute difference between two empirical cumulative distribution functions.
χ^2	**McNemar Test Statistic**	Dimensionless	McNemar test statistic, the squared difference of the two discordant counts divided by their sum.
$\mathcal{S}$	**Mechanism Output Set**	Set	Set of possible outputs of a randomized mechanism over which a differential privacy bound must hold.
δ_{MDE}	**Minimum Detectable Effect**	Dimensionless	Smallest absolute difference in the outcome metric that an A/B test is designed to detect.
x_{feat}	**Monitored Feature Level**	Varies	One value of a monitored feature, the argument at which its distribution or empirical cumulative distribution function is evaluated.
I	**Mutual Information**	Function	Mutual information between two random variables, how much observing one reduces uncertainty about the other.

Symbol	Definition	Unit	Notes
X	**Neighboring Dataset**	Record	Dataset input to a randomized mechanism in the differential privacy definition, compared against a copy that differs in one record.
O	**Observable Model Output**	Record	Model output an adversary can observe, such as a prediction, score or generated text.
$p_{\mathbf{val}}$	**Observed Significance Level**	Dimensionless	P-value of a hypothesis test, the probability under the null hypothesis of a result at least as extreme as the one observed.
$\mathbf{s}$	**Pairwise Mask Secret**	Vector	Shared random mask vector two clients derive from a common secret, which one adds and the other subtracts so that it cancels in the secure-aggregation sum.
$n_{\mathbf{bins}}$	**Population Stability Bin Count**	Integer	Number of fixed bins over which the population stability index sums.
z_β	**Power Critical Quantile**	Dimensionless	Positive standard normal critical value corresponding to the chosen statistical power of a test.
$\hat{Y}$	**Predicted Decision**	Dimensionless	Binary decision a classifier outputs for a randomly drawn individual, one for a positive prediction.
σ	**Privacy Noise Multiplier**	Dimensionless	Ratio of the Gaussian noise standard deviation to the gradient clipping norm in differentially private training.
z	**Proportion Test Statistic**	Dimensionless	Test statistic of a two-proportion z-test, the difference of two sample proportions divided by its pooled standard error.
G	**Protected Group Membership**	Record	Protected group a randomly drawn individual belongs to, as a random variable over group labels.
R	**Protected Training Record**	Record	Training record of one protected individual, the unit whose leakage a privacy claim bounds.
$\mathcal{M}$	**Randomized Mechanism**	Function	Randomized algorithm that maps a dataset to a released output, the object a differential privacy guarantee constrains.
b	**Reference Bin Index**	Integer	Position of one reference bin among the bins a continuous feature is discretized into for the population stability index.
E	**Reference Bin Proportion**	Dimensionless	Proportion of reference (training-time) observations falling in one bin of the population stability index.
$\alpha_{\mathbf{R}}$	**Rényi Order**	Dimensionless	Order of the Rényi divergence at which a Rényi differential privacy guarantee is stated, a number greater than one.
n	**Required Users Per Variant**	Integer	Number of users each variant of an A/B test needs to detect the minimum detectable effect at the chosen confidence and power.
b	**Second Compared Group Label**	Record	Second of two group labels compared in a group-fairness condition, one value of the protected group attribute.
H	**Statistical Hypothesis**	Record	A statistical hypothesis about a population, such as the null hypothesis or the alternative hypothesis.
g	**Subgroup Index**	Integer	Label of one population subgroup whose outcomes are compared against a reference group.
P_0	**Training Distribution**	Distribution	Data distribution the model was trained on, the baseline that drift is measured against.
Y	**True Outcome**	Dimensionless	True binary outcome of a randomly drawn individual, such as whether an applicant repays, one for the positive class.
v	**True Outcome Level**	Integer	One value of a binary true outcome, zero or one, on which a fairness condition is evaluated.

Distributed governance, reliability, and safety constants

A few thresholds and limits recur across the governance and reliability chapters, from classifier decision thresholds to failure-detection timeouts and sequential test statistics. The table below collects these shared constants so their symbols stay distinct from model parameters and network latencies:

Symbol	Definition	Unit	Notes
$h(x)$	**Classifier**	Function	Binary classifier $h(x)$ whose positive-prediction rates the fairness criteria constrain.
C	**Gradient Clipping Norm**	Scalar	Per-example gradient clipping bound C in differentially private training; not a C^3 axis.
T_{deadline}	**Deadline**	seconds	Target latency deadline for batch or phase completion.
T_{timeout}	**Failure-Detection Timeout**	seconds	Interval without a heartbeat after which the coordinator declares a rank failed.
τ	**Response Threshold**	Scalar	Decision or classification threshold τ.

Units, Precision, and Disambiguation

A link rate quoted in gigabits per second looks eight times larger than the same rate in gigabytes per second, and an 80 GB memory label can mean either 80×10^9 or 80×2^{30} bytes depending on whether decimal or binary prefixes apply. The unit, precision, and disambiguation conventions below keep numerical comparisons consistent across every chapter.

Units and precision

- Physical units: This book uses SI (metric) units throughout, including meters, kilograms, seconds, watts, and °C, consistent with standard engineering and scientific practice. Where source data was originally reported in imperial units, the book converts to SI and notes the original values parenthetically. A space always separates the number from the unit (for example, 100 ms, 2 TB/s).
- Data and memory: This book uses **decimal SI prefixes only**: KB = 10^3 bytes, MB = 10^6, GB = 10^9, TB = 10^{12}. Binary-prefixed units do not appear in prose; model sizes, activation footprints, throughputs, and transfer volumes are reported in decimal units (for example, 2 TB/s, 102 MB). Hardware memory capacities retain the vendor's nameplate figure and label. A modeled capacity of 80×2^{30} bytes follows the registered device assumption, not a guarantee implied by an 80 GB label. Where a worked example subtracts a computed footprint from a device capacity, the registered capacity in bytes and the footprint are both converted to decimal; Appendix G states this exception in full.
- Compute: The notation distinguishes operation counts from rates. Total work uses FLOPs (for example, GFLOPs, TFLOPs), while throughput uses FLOP/s with decimal prefixes (for example, GFLOP/s, TFLOP/s).
  - 1 TFLOP = 10^{12} FLOPs
  - 1 TFLOP/s = 10^{12} FLOPs per second
  - Arithmetic intensity conventionally uses FLOP/byte as a unit ratio (floating-point operations per byte moved).
- Network throughput: Reported in bytes/s (GB/s, TB/s) or bits/s (Gbps) depending on convention. InfiniBand and Ethernet specifications use Gbps; application-level throughput uses GB/s. The text states the convention on first use.
- Power: Power carries a subscript (P_{avg}, P_{device}, P_{cluster}, P_{IT}); bare P remains parameter count.
- Latency percentiles: Written p50, p95, and p99 in upright type, inside math as well as in prose.
- Currency: Dollar amounts use the dollar sign (\$); unless otherwise noted, dollar-denominated costs are U.S. dollars (USD).
- Precision:
  - **FP64**: Double precision (8 bytes)
  - **FP32**: Single precision (4 bytes)
  - **TF32**: TensorFloat-32 (19-bit Tensor Core compute mode; inputs and storage remain FP32)
  - **FP16**: Half precision (2 bytes, standard range)
  - **BF16**: Brain float (2 bytes, wide dynamic range)

- **FP8**: Quarter precision (1 byte, E4M3 or E5M2 format)
- **FP4**: 4-bit floating-point format
- **INT8**: 8-bit integer (1 byte)
- **INT4**: 4-bit integer; lower integer precisions follow the same uppercase `INTn` pattern (for example, INT3, INT2)

Quick reference: Resolving collisions

Several single-letter symbols carry conflicting default meanings in machine learning and systems papers. The table below resolves the four primary single-device collisions:

Symbol	ML meaning	Systems meaning	Book convention
B	**Batch Size**	Bandwidth	**Batch Size**. Use BW for bandwidth.
D	**Dataset Size**	Data Volume	**Dataset Size**. Use D_{vol} for bytes moved.
η	**Learning Rate**	Efficiency	**Learning Rate**. Use η_{hw} for efficiency.
P	**Parameters**	Peak FLOP/s	**Parameters**. Use R_{peak} for peak rate.

As a general principle, ML conventions take precedence for single letters, while systems concepts get subscripts or multi-letter symbols. This reflects the primary audience (ML practitioners learning systems) and preserves compatibility with the ML literature.

Additional collision points in distributed execution

Moving from a single accelerator to a multi-node cluster introduces symbol collisions across networking, queuing theory, and fault tolerance beyond those of single-device notation:

Symbol	ML meaning	Distributed meaning	Book convention
N	(none)	Number of Devices	**Device Count** in distributed context (accelerators, not host nodes); $N = N_{\text{GPU}}$.
M	GEMM dimension; memory root (M_{total})	Collective Payload	**Payload** when bare, outside GEMM shape statements. Memory footprints always carry a subscript.
λ	Regularization weight	Failure Rate	**Subscript every other use**. Bare λ is the per-component failure rate; arrival rate is λ_{arr}.
d, t, p	Hidden dimension; time; distribution $p(\cdot)$	Data-, tensor-, and pipeline-parallel degree	**Core meanings**. Parallelism degrees are N_{DP}, N_{TP}, and N_{PP}.
τ	Temperature or threshold	Checkpoint Interval or Delay	**Subscript every distributed use**. Checkpoint intervals are T_{ckpt} and T^*_{ckpt}; gradient staleness is τ_{stale}.
α	Significance Level	Network Latency	**Subscript every other use**. Bare α is α-β per-message latency; α_{scale} is the scaling exponent, α_{sig} the significance level, and α_{R} the Rényi order.
ρ	Correlation	Utilization or Ratio	**Subscript every other use**. Bare ρ is the communication-computation ratio; ρ_{serv} is serving or queue utilization.
ϵ	Small Constant	Privacy Budget	**Subscript every other use**. Bare ϵ is the privacy budget; ϵ_{adv} is the adversarial radius, and stabilizing constants carry their own subscripts.
g, G	Gradient $\mathbf{g}$	LogGP gaps; GPUs per node	**LogGP only** when bare. Gradients are bold $\mathbf{g}$; accelerators per node are G_{node}.
K	Top-K, key matrix $\mathbf{K}$	Attack Steps	**Local when bare**. PGD attack steps are K_{PGD}; other counts are defined at use or subscripted (K_{top}).

MAIN MATTER

1

Introduction

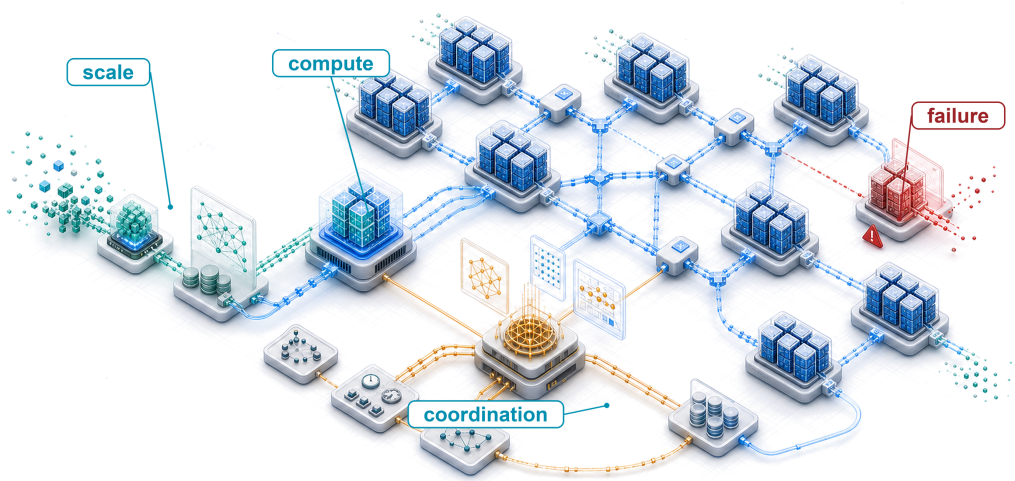

1.1 The Scale Moment
1.2 From Single Node to Fleet
1.3 AI Scaling Laws
1.4 Constraints of Scale
1.5 The C^3 Taxonomy and Foundations of Scale
1.6 The Fleet Stack and the Lighthouse Archetypes
1.7 The Structure of This Textbook
1.8 Fallacies and Pitfalls
1.9 Summary

Purpose

Why are single-machine engineering principles insufficient for machine learning at fleet scale?

Machine learning at scale has a physics of its own. On one node, performance is often governed by the memory wall: data must reach the accelerator fast enough to keep the math busy. Scaling laws explain why one node stops being enough, because model quality improves predictably but sublinearly with compute, data, and parameters, so each further gain demands a larger fleet. In a distributed cluster, the memory problem stretches across machines. Data must move through network links between racks, and the slowest shared crossing can limit the entire job. Hardware failures also change character. A failure that stops one local training run can stall a whole distributed job when thousands of workers must advance together. This discontinuity explains why mastery of single-machine ML is no longer sufficient. A faster device may finish its local work sooner yet spend the saved time waiting for other machines. Adding devices can increase capacity while creating more shared work to coordinate and more components to recover. Serving introduces another consequence: spare capacity is useful only if requests can reach it within their deadlines. The engineering decision must therefore account for the behavior of the whole fleet. That wider boundary also amplifies societal impact because these systems may serve billions of users. Bias in any model can cause harm, but bias in a widely reused foundation model can propagate across many downstream systems. In C^3 terms, compute, communication, and coordination turn algorithmic choices into questions of physical topology, fault tolerance, security, and governance.

Part IV
Governance
Assurance
Part III
Operations
Serving
Part II
Control
Execution
Part I
Fabric
Facility

Learning Objectives

- Explain why leaving one machine shifts the binding constraint from local memory to the network, energy, and routine failure
- Explain how empirical scaling laws drive fleet growth and where breakdown conditions bound compute-optimal allocation
- Diagnose fleet-scale distribution risks using synchronous tight coupling, reliability gaps, consistency-availability trade-offs, and communication intensity
- Explain how Compute, Communication, and Coordination extend the Data-Algorithm-Machine lens for production fleets
- Apply the Fleet Law to separate a training step's useful compute from its communication and coordination overhead
- Explain how the fleet stack connects infrastructure decisions to deployment and governance constraints

1.1 The Scale Moment

Machine learning on a single accelerator is governed by local memory hierarchy and arithmetic intensity, the physics of silicon. The boundary of single-machine systems appears when that local machine is no longer the system. As models move from research prototypes to global services, the binding constraints move outward into racks, networks, power delivery, and recovery machinery. That change is not a smooth extrapolation; it is the **scale moment**, the physical and operational transformation that occurs when models encounter three recurring walls: memory, network, and energy.

On one accelerator the memory wall often binds; after distributing a communication-intensive workload across thousands of nodes, the network wall can replace local memory as the primary performance bottleneck. Its hard limits are speed-of-light latency and bisection bandwidth, the aggregate bandwidth across the narrowest cut that splits the fleet into two equal halves. On a given fabric, adding workers does not raise that bound, so a tightly coupled job whose traffic must cross the cut stops gaining speed once the cut saturates. Serving billions of users around the clock can then encounter the energy wall, the limit on how much power a facility can deliver and remove as heat, making thermodynamic efficiency a first-order engineering requirement. Crossing those walls also opens a reliability gap, making hardware failure a routine event rather than a rare exception.

A synchronous distributed training step gives this change a concrete form. Each worker computes part of the update, exchanges model information with other workers, and waits until the required participants can advance together. These activities give the book its C^3 lens: Compute is local execution, Communication moves data between workers, and Coordination keeps their progress and state consistent, including after failures. The Fleet Law expresses the corresponding time budget. Dividing a fixed workload across more devices can shorten local computation, but the step also pays for communication and coordination, less whatever work can overlap. Faster accelerators help only when the remaining costs allow the saved compute time to shorten the whole step.

This book assumes familiarity with single-machine training, inference, and accelerator memory. Readers who want to refresh that foundation should begin with Appendix A, which restates it in reference form. The appendix covers the Data · Algorithm · Machine (D·A·M) lens that organizes single-machine diagnosis, the iron law, the Roofline Model and its arithmetic-intensity boundary, the accelerator memory hierarchy, the memory cost of training state per parameter, and the arithmetic of transformer layers.

The book's aim is to develop engineering judgment across infrastructure, distributed execution, deployment, and governance, so that a local improvement can be evaluated against the behavior of the whole fleet. This chapter builds that judgment in the order a fleet engineer meets the problem. It first traces how far training compute and cluster size have grown and what changes when a job leaves one machine. It then turns to the empirical scaling laws that keep driving that growth, the physical, logical, and societal constraints that distribution imposes, and the C^3 taxonomy and Fleet Law that diagnose where a distributed step loses time. It closes with the fleet stack, an eight-row

layered map of the fleet that locates each cost in the architecture, and three recurring workloads that test how the dominant cost changes. Parallelism, communication, and recovery techniques appear here only by name. The task in this chapter is to identify the constraint each one must address, and later chapters develop the mechanisms.

Between 2012 and the mid-2020s, public and third-party estimates place training compute growth from roughly 10^{18} FLOPs for AlexNet to approaching 10^{25} FLOPs for leading large models, seven to eight orders of magnitude. The difference is qualitative, not merely quantitative. A workload that once fit on two GPUs now spans tens of thousands of accelerators. Recent capability leaps have repeatedly come from applying computation at previously impractical scales, and the empirical scaling laws examined in section 1.3 explain why organizations keep paying for each further order of magnitude.

This growth unfolded across a sequence of landmark models. AlexNet trained on two GTX 580 GPUs for approximately 5–6 days (Krizhevsky et al. 2012). BERT required 64 Tensor Processing Unit (TPU)[1] chips for 4 days, roughly 6,144 chip hours (Devlin et al. 2019).

GPT-3 consumed an estimated 3.14×10^{23} FLOPs during training (Brown et al. 2020). A single A100, a generation newer than the V100 GPUs that trained it, would need about 32 years to perform that work even at its full 312 TFLOP/s peak, which is the plainest answer to why one node is no longer enough. A later carbon accounting of the run itself pairs the budget with 10,000 V100 GPUs sustaining about 24.6 TFLOP/s each (roughly 20 percent of the V100's peak arithmetic rate) for roughly 15 days of training (Patterson et al. 2021). The gap between that delivered rate and the datasheet peak raises a question this book keeps returning to, namely how much of each accelerator a fleet actually puts to use.

PaLM trained on 6,144 TPU v4 chips for 50 days and then on 3,072 for 14 days more, consuming approximately 10^{24} FLOPs (Chowdhery et al. 2022). OpenAI's GPT-4 report did not disclose model size, training hardware, or training compute, so the GPT-4-class scenario in this section should be read as illustrative rather than a published configuration (OpenAI et al. 2023). These examples sit within the rapid training-compute growth trend documented by earlier compute-trend studies (Amodei and Hernandez 2018; Sevilla et al. 2022). Together, these landmarks show training budgets moving from single-machine experiments into distributed systems where hardware count, wall-clock time, and operational coordination are part of the model design (table 1.1).

[1] **Tensor processing unit (TPU)**: Google's custom application-specific integrated circuit (ASIC) for machine learning, whose matrix units are systolic arrays, grids of multiply-accumulate cells that pass data between neighboring cells (Jouppi et al. 2017). At fleet scale, the inter-chip network matters as much as the array: TPU v4 joins its chips with an optically reconfigurable interconnect, so the topology of the network, not only the speed of each chip, determines how well many accelerators cooperate (Jouppi et al. 2023).

Table 1.1: Training Compute Evolution: Growth in hardware scale and training duration across published landmarks and an illustrative GPT-4-class scenario, showing why distributed systems become central as training budgets move from 10^{18}-scale to the largest training runs.

Model	Year	GPUs/TPUs	Training time	Estimated FLOPs
AlexNet	2012	2 GPUs	5–6 days	~10^{18}
BERT-Large	2018	64 TPUs	4 days	~10^{20}
GPT-3	2020	10,000 V100 GPUs (run est.)	~15 days	3.14×10^{23} FLOPs
PaLM	2022	6,144 TPUs	~64 days	~10^{24}
GPT-4-class scenario	2023	~25,000 GPUs (illustrative)	~100 days (illustrative)	~10^{25} scenario

Epoch AI's database of notable models records, for each training run it can document, the compute the run consumed together with how many accelerators it used and for how long (Epoch AI 2026c). Plotted run by run, training compute climbs steadily on a logarithmic scale, and it separates early single-node vision models from multi-node transformer training (figure 1.1, panel a). Because training compute is the product of the accelerator count, the training duration, and the rate each accelerator delivers, the same records also show where the growth came from. Across the 146 runs from 2012 through 2024 that report all three, compute grew about 3.9× per year. The accelerator count grew 1.6× per year, the duration 1.3×, and the delivered rate per accelerator, which folds in both faster hardware and better utilization, 2× (figure 1.1, panel b). A higher delivered rate per

accelerator supplied roughly half of that growth. More accelerators running for longer supplied the other 51 percent, and that share must come from the fleet rather than from any single accelerator.

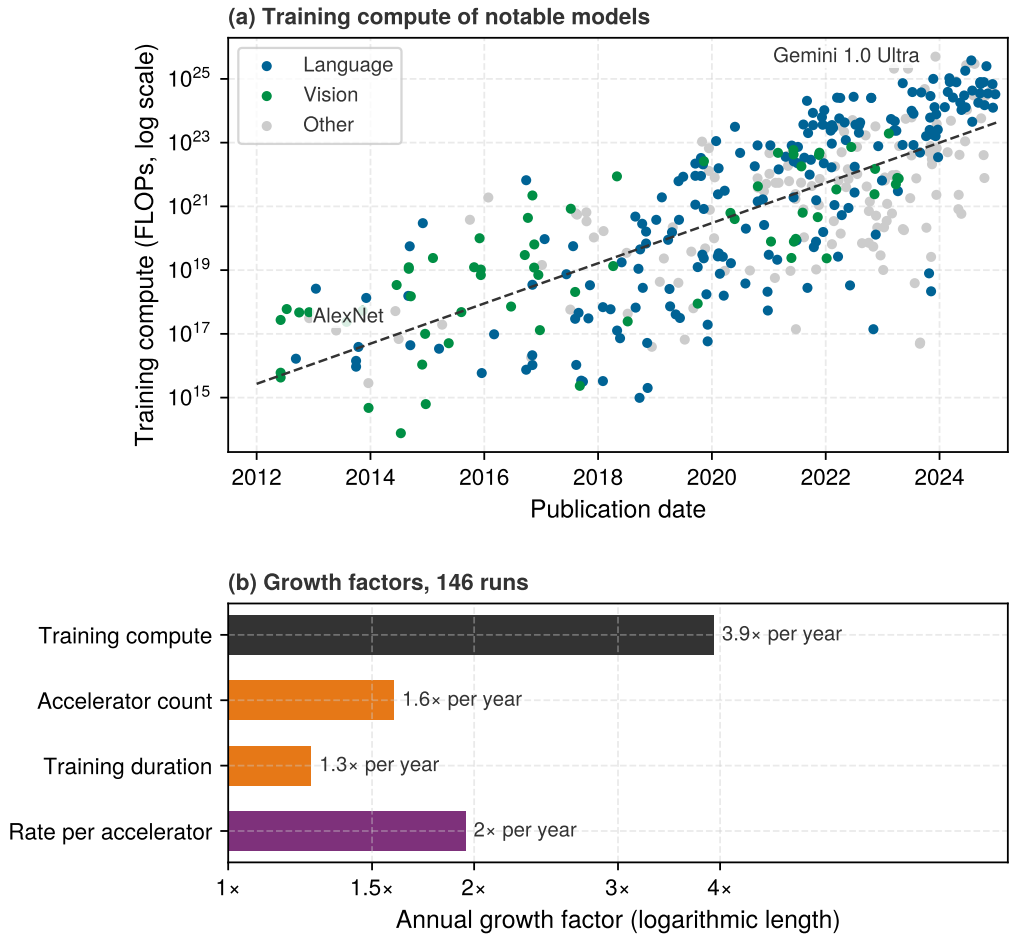

Figure 1.1: Where Training-Compute Growth Comes From: (a) Training compute of notable models published from 2012 through 2024 on a logarithmic axis, colored by domain, spanning about eight orders of magnitude from AlexNet to the largest runs; the dashed line is a least-squares fit. (b) Annual growth factors fitted to the runs that report accelerator count and training duration. Compute growth is the product of the other three bars, and the bar lengths are drawn on a logarithmic scale so each factor's share of the growth reads directly as its length. The delivered rate per accelerator includes both hardware speed and utilization. Data from Epoch AI (2026c); many values are Epoch AI estimates.

The language-model points pull away from the vision points because compute demands escalated sharply as transformer architectures replaced convolutional pipelines. Landmark frontier runs grew their accelerator counts faster than the average run, roughly doubling every year (figure 1.2), a rate that quickly outpaced what any single node could host.

Across the four-order-of-magnitude climb charted in figure 1.2, accelerator counts rise from two GPUs for AlexNet to 16,384 H100 GPUs for Llama 3 along an annual doubling trajectory ($\sim 2\times$ per year). The one estimated point is GPT-3's, placed at the size of the Microsoft supercomputer built for OpenAI in 2020, more than 10,000 GPUs (Langston 2020). This empirical trajectory is the foundation of the scale moment, because leading models have required far larger accelerator fleets than earlier deep-learning landmarks. The curve alone, however, does not explain which constraint breaks first. Large language-model training, recommendation serving, and federated mobile learning all encounter fleet scale, but each stresses a different part of the system.

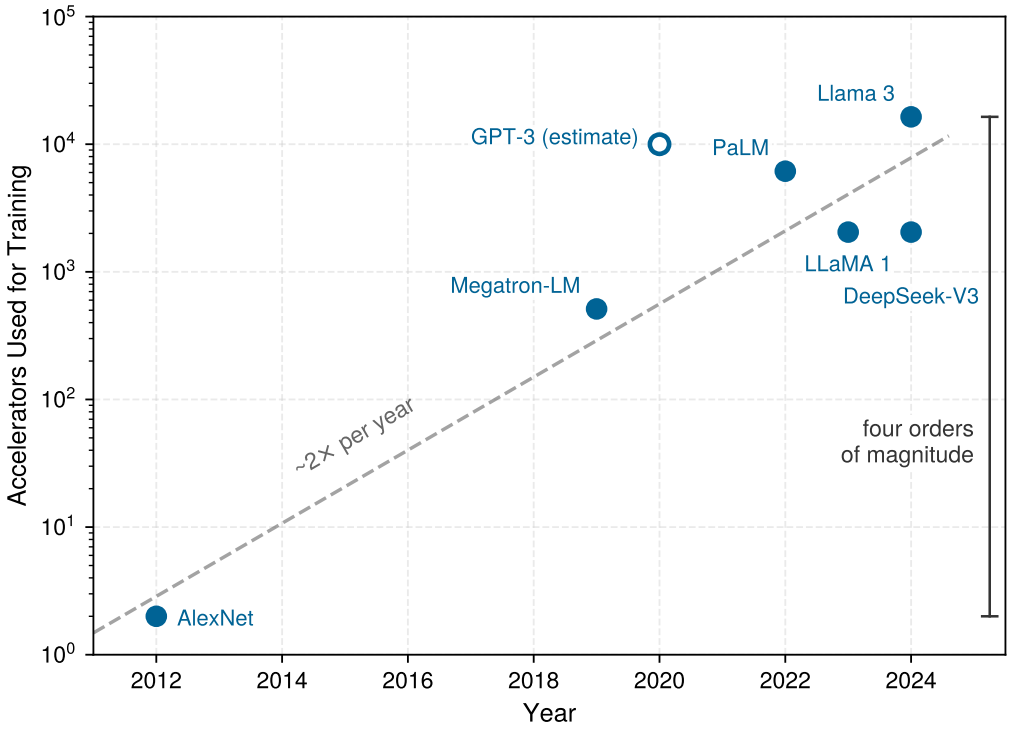

Figure 1.2: The Cluster Size Explosion: Number of accelerators used to train landmark models, 2012–2024. Verified counts from published papers are shown as filled circles; the GPT-3 estimate (hollow marker) reflects approximate cluster size from Microsoft infrastructure announcements rather than a precise published count. The dashed trend line indicates approximately 2× annual growth in cluster size, a rate that outpaces Moore's law and motivates the infrastructure challenges developed throughout the book.

Lighthouse 1.1: Lighthouse archetypes at scale

Lighthouse archetypes are the canonical workloads that keep the scale moment concrete throughout this book. Their role at this point is to show that "more machines" is not one engineering problem, because the dominant constraint depends on what must cross the fleet boundary. The full roster with the C^3 taxonomy mapping and binding-constraint analysis appears at section 1.6.1.

These archetypes turn the growth curve into an engineering question. The same increase in accelerator count can produce a network wall for a large language model split across devices, either within each layer (tensor parallelism) or by assigning consecutive groups of layers to different devices as pipeline stages (pipeline parallelism). It can produce a coordination bottleneck for a recommendation model whose embedding tables, which map each user and item to a vector, are spread across hundreds of nodes, so each request must route its sparse lookups to the many nodes that hold its rows. It can also produce a per-device compute ceiling for federated learning, which trains on millions of phones without collecting their raw data and so takes the problem beyond the data center.

↪ Data, tensor, and pipeline parallelism, and the traffic each places on the fabric, are developed in chapter 6.

Moving from one machine to many changes the governing physics of the system, not merely its device count. A single node is bound locally by the memory wall between its accelerators and their high-bandwidth memory, and it has no peers to wait for. A fleet executes across a shared network fabric under a scheduler that places and recovers the job, which exposes the workload to the network wall, to stragglers, and to routine node failures (figure 1.3).

↪ Federated rounds across millions of devices, where dropouts, stragglers, and per-device budgets bind, are developed in section 11.7.

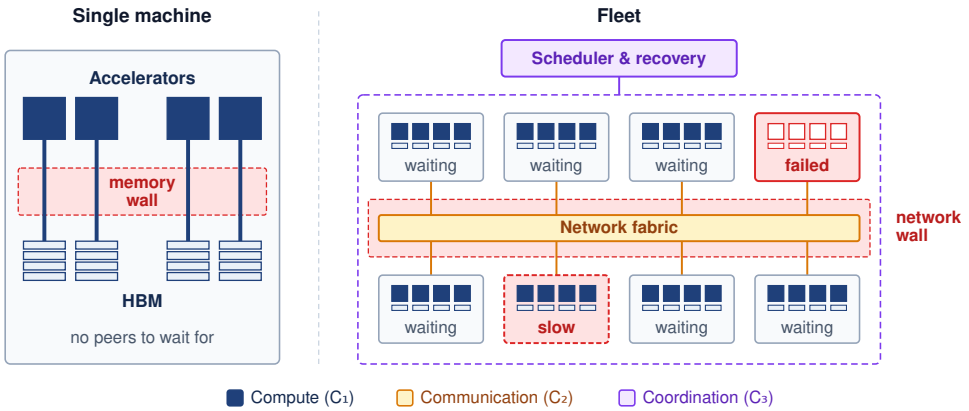

Figure 1.3: The Scale Discontinuity: On a single machine (left), the usual binding constraint is the memory wall between the accelerators and their high-bandwidth memory, and there are no peers to wait for. In a fleet (right), nodes exchange state over a shared network fabric under a scheduler that places the job and recovers it, so the three C^3 roles appear as separate components. The fabric can become the binding network wall, and failed and slow nodes become routine events that hold back the whole job.

At fleet scale, the binding constraint may shift from local silicon to the network fabric, while resilience joins optimization as a primary concern. The unit of compute is no longer a single server but a **machine learning fleet**, an interconnected distributed system designed to act as a coherent engine.

> **Definition 1.1: Machine learning fleet**
>
> ***Machine learning fleet*** is a distributed collection of compute devices, storage, and networks managed together to train or serve machine learning models.
>
> 1. **Significance**: In tightly coupled synchronous training, the total time T can be governed by the slowest worker (straggler). Sustaining that workload requires balancing bisection bandwidth ($\text{BW}_{\text{bisect}}$) against aggregate compute capacity (R_{peak}). Serving fleets and heterogeneous edge fleets impose different communication and coordination requirements.
> 2. **Distinction**: The fleet is the unit of engineering, but its workers need not all advance together. Synchronous tight coupling makes a training job depend on a shared step barrier; serving replicas can process separate requests, and federated devices can participate intermittently. The workload determines the required coupling and recovery behavior.
> 3. **Common pitfall**: Treating an ML fleet as "just more servers" overlooks shared network, placement, and recovery constraints. A tightly coupled data-center fleet behaves as one machine, the warehouse-scale computer of section 1.2.2; an edge fleet extends the engineering boundary to devices outside that managed facility.

At the scale of thousands of GPUs, failure modes appear that a single GPU never encounters. A network partition (a fault that splits the fleet into groups that cannot reach one another) can leave a synchronous job hung at a collective (an operation, such as a gradient exchange, that every worker must join) until a timeout aborts it, and in loosely synchronized designs it can let groups of workers drift apart. Stragglers (workers that process data slower than peers due to hardware variation or thermal throttling) can bottleneck entire training runs. Even infrequent component failures accumulate across a large fleet, so long training jobs must expect interruptions.[2] Systems must preserve recoverable state frequently enough to bound the work lost at each interruption.

The interruption cadence follows from simple arithmetic. Consider a hypothetical GPT-4-class training run on A100-class hardware, a cluster of 25,000 GPUs running for 100 days, values that follow a third-party estimate (Patel and Wong 2023) rather than an OpenAI disclosure (OpenAI et

[2] **Hardware failure rates at scale**: The arithmetic below uses an assumed per-GPU mean time to failure and independent failures to illustrate how interruption frequency grows with fleet size. Production job interruptions include software, networking, and maintenance as well as accelerator faults. Meta's Llama 3 report provides an observed example of this broader interruption category (Dubey et al. 2024); it does not establish a universal annual GPU failure rate.

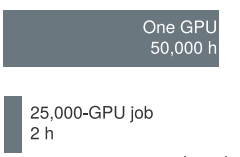

A 25,000-GPU job fails 25,000 times as often as one GPU.

al. 2023). With an assumed per-GPU mean time to failure (MTTF) of 50,000 hours and independent failures, the system mean time between failures (MTBF) is the per-GPU MTTF divided by the GPU count, about 2 hours, or roughly 12 interruptions per day.

Production runs show a similar cadence. Meta's Llama 3 405B pretraining on 16,384 H100 GPUs recorded 419 unexpected interruptions over a 54-day snapshot, about one every 3 hours, from software, network, and host faults as well as GPU faults (Dubey et al. 2024). At that rate, a job that saved its state only every few hours would often be interrupted before its next save, and manual restart cannot keep pace, so recovery must be automated and designed into the system from the outset (section 7.2.2 follows how this run recovered).

↪ The failure-rate arithmetic that turns a per-device MTTF into a job's interruption cadence, checked against production failure records, is developed in section 7.2.

Organizations answered these challenges with purpose-built fleets whose defining component is the network as much as the accelerator. Meta's Research SuperCluster, announced in 2022, was designed to reach 16,000 NVIDIA A100 GPUs connected by 200 Gb/s InfiniBand[3] links (AI 2022). Google's TPU v4 pods join 4,096 chips delivering 1.1 EFLOP/s of aggregate peak compute (Zu et al. 2024). Later Azure AI data center announcements describe training fabrics of tens of thousands of GPUs (Microsoft 2025). Each accelerator's share of that aggregate compute must exchange its results through links of a few hundred gigabits per second.

The scale moment establishes why exponential growth in compute demand forces some ML systems beyond any single machine, making the network fabric a potential bottleneck and component failure routine over long runs. The next question is how to organize the engineering response. Established distributed-systems techniques remain essential, but global barriers and large shared model state change their costs and determine which recovery and consistency mechanisms apply. Tracing those changes means following how the layers of one machine map onto a fleet, beginning with how sharply the links between servers fall below the bandwidth inside one.

3 **InfiniBand**: A data-center interconnect whose defining feature for ML is remote direct memory access (RDMA), which bypasses the OS kernel to move data directly between application memory on different machines at microsecond-scale latency. HDR InfiniBand provides about 25 GB/s line-rate bandwidth per link and NDR about 50 GB/s; that per-link bandwidth, relative to the computation each worker performs between exchanges, decides whether large-model training is compute-bound or communication-bound. RDMA-capable Ethernet (RoCE) now offers comparable line rates.

1.2 From Single Node to Fleet

Inside one H100 server, NVLink moves tensors between accelerators at 450 GB/s in each direction per accelerator. Once a training job spans racks, the same tensors leave through an InfiniBand NDR port at 400 Gb/s (50 GB/s) in each direction, a bandwidth step down of roughly 9× for an accelerator served by one such port. That step illustrates a general pattern. The layered structure of an ML system persists across scale, while the physics at each layer changes.

The **fleet stack** names the fleet's layers by grouping them into four bands, Infrastructure at the bottom, then Distribution, Serving, and Governance at the top, with two rows in each (figure 1.4). At the bottom, the hardware substrate of one node grows into the Infrastructure band, whose Facility and Accelerators row houses many nodes and whose Fabric and Storage row joins them with an RDMA fabric spanning racks. There the binding constraint can shift from the memory wall to the network wall as network congestion and message-passing latency consume a growing share of step time. Efficient data movement inside each node remains a prerequisite, because a fleet of underutilized nodes only multiplies the waste. One layer up, system software (a single CUDA runtime managing PCIe direct memory access (DMA)) becomes the Distribution band. This band holds the NVIDIA Collective Communications Library (NCCL) and RDMA libraries that coordinate thousands of processes across the fabric, together with the schedulers that place distributed jobs and the recovery machinery that restarts them.

The upper bands add responsibilities rather than replace the familiar software layers. The ML framework (PyTorch or JAX executing a training loop on one node) remains part of model execution. The Serving band adds request handling and deployment pipelines that manage rolling releases across the fleet. Applications also remain, while the Governance band sets security, robustness, sustainability, and responsible AI requirements for their operation. The full stack and its eight rows, the map that organizes the rest of this book, return in section 1.6. The stack identifies where a design decision belongs; C^3 identifies how that decision affects execution time.

A map of layers does not by itself explain why the same decision behaves differently once it spans a fleet. Four changes separate a fleet from a larger server: synchronous coupling turns independent tasks into a shared step, the facility itself becomes the computer, communication can dominate the step, and failure becomes routine. The scale moment quantified that last change, and its availability model follows with the other constraints of scale in section 1.4.1. Archetype A (GPT-4/Llama-3), the

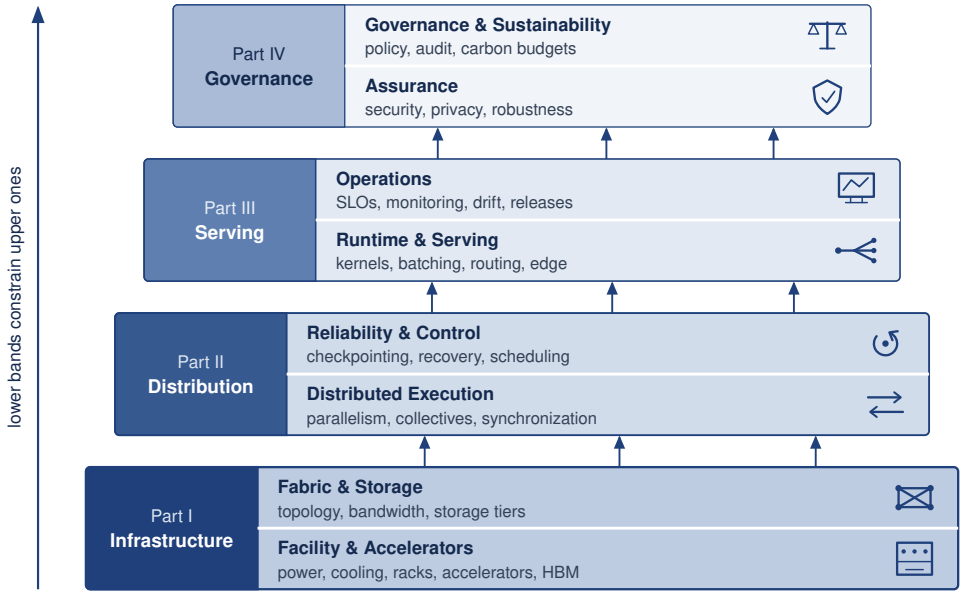

Figure 1.4: The Fleet Stack: The organizing framework for this book, eight rows grouped into four bands, one band per part. Each band is labeled with its part number and a short functional name for that part, and each row lists what it holds beside the icon that marks it in the chapter margin stacks; the Reliability & Control row, for example, is where orchestration, scheduling, and recovery live. Capabilities decided in a lower band constrain the choices available in the bands above it. The same stack, shaded to show each chapter's focus, appears beside the Purpose of every chapter.

large-language-model workload of the book's canonical roster (section 1.6.1), runs on a reference cluster of NVIDIA H100 accelerators and makes each change concrete where a chapter needs it.

1.2.1 Independent tasks vs. synchronous training steps

When one worker in a MapReduce batch job fails or slows down, the scheduler can retry its task alone while the rest of the cluster keeps making progress. When one worker in a synchronous training job slows by 10 percent on the critical path (the chain of dependent work that sets the step's length), every other accelerator waits at the step barrier and the whole job slows by nearly 10 percent. Many distributed systems isolate work into tasks or requests so that one failure need not stop unrelated work: **MapReduce** (Dean and Ghemawat 2004) partitions an input into map tasks and combines their outputs in reduce tasks, both of which can be retried; Apache Spark (Zaharia et al. 2016) processes data in parallel stages that exchange results at explicit shuffle boundaries; and a web microservice handles requests whose failures can be isolated. Search engines and banking databases use other forms of replication, partitioning, and coordination; they are not uniformly independent or asynchronous. The useful contrast is therefore between failure domains (what a single fault can stop) and coupling patterns, not between "traditional" and ML systems.

Large synchronous training jobs operate under **synchronous tight coupling**, in which all required workers must finish one shared step before any of them can begin the next. Parameter-server architectures (Li et al. 2014), which keep the model on dedicated servers, and collective runtimes manage distributed state, but synchronous optimization updates a large shared model state across thousands of accelerators at every step rather than completing independent one-shot tasks. At each synchronization point, progress can depend on the slowest participating worker, and communication-heavy partitions can also become bisection-bandwidth bound because large gradient or activation tensors cross the network every step. These properties do not invalidate established recovery techniques; they increase recovery cost and restrict which replication, re-execution, and consistency strategies preserve useful training progress.

↪ The barrier that makes every update wait for every required contribution, and the failure modes it creates, are developed in section 6.3.1.6.

The two execution pathways differ at the point of failure (figure 1.5). In MapReduce, a failed map task is retried on its own while its peers complete. In synchronous data-parallel training, every worker holds a full copy of the model, spread over as many accelerators as its training state needs, and processes a different slice of each batch. Before the shared parameter update can occur, every worker must contribute its gradient to an AllReduce, a collective operation that combines the gradients of all workers and returns the result to each of them. A slow map task can delay downstream reduction and job completion, but other map tasks finish without waiting; in synchronous training, by contrast, healthy workers sit idle whenever a straggler delays the collective.

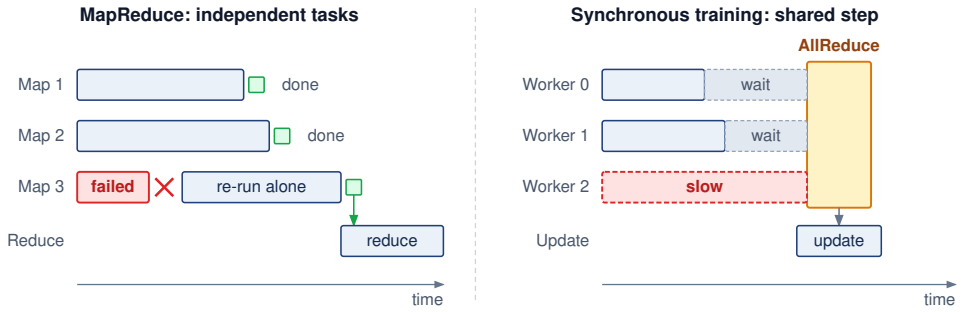

Figure 1.5: Retryable Tasks vs. Shared Training Steps: Map tasks (left) deliver their outputs independently, so a failed task is re-run on its own while the others finish; its missing output delays only the reducers that need it. In synchronous data-parallel training (right), the AllReduce needs every worker's gradient before the shared update can proceed, so a slow worker on the critical path holds the whole step and the healthy workers wait. How much the step stretches depends on overlap and the other costs of the step.

The collective barrier in figure 1.5 is why MapReduce-style localized task re-execution cannot be transferred directly to tightly coupled training. When all workers advance a single shared parameter step, recovery must restore mutually consistent model, optimizer, and input state across every participating node before coordinated execution can resume.

↪ Coordinated checkpointing, which captures that mutually consistent state so that recovery need not search for it, is developed in section 7.8.5.

Checkpoint 1.1: The fleet mindset

These questions check whether the fleet mindset is clear:

- ☐ **Straggler impact**: compare how a slow map task can delay job completion with how a persistent straggler can delay repeated synchronous training steps.
- ☐ **Bisection bottleneck**: explain when frequent cross-rack tensor exchanges would make synchronous training more bandwidth-sensitive than a service handling mostly independent requests.
- ☐ **Tight coupling**: explain synchronous tight coupling and connect it to the need for all participating contributions before a shared training update completes.
- ☐ **Linear speedup myth**: using the barrier argument above, explain how a straggler can keep a synchronous job with twice the GPUs from halving its step time.

1.2.2 The shift to the warehouse-scale computer

Once a single training job spans thousands of accelerators across dozens of racks, a facility's cooling loops, power distribution, and bisection bandwidth can cap step throughput long before any individual chip reaches its arithmetic limit. A tightly coupled data-center fleet therefore demands the **warehouse-scale computer (WSC)**[4] perspective. Where a machine-centric view treats the data center as a building that houses many separate servers, the warehouse-scale view treats the data center itself as one computer. Single-node abstractions then map directly onto distributed infrastructure: the network fabric functions as the system bus, distributed storage serves as the local disk, and the fleet orchestrator acts as the operating system.

[4] **Warehouse-scale computer (WSC)**: Barroso et al. (2019) frame the data center itself as the computer. For ML fleets, power delivery, cooling topology, optical network layout, and bisection bandwidth become first-order constraints; the power limit of a 100 MW facility can cap training throughput before any single accelerator's TFLOP/s does.

For this data-center regime, the engineering boundary extends from local execution to the shared power, network, and recovery systems that sustain the entire facility. The warehouse-scale computer, the standard design for large models, moves data between chips over a network. Wafer-scale engines, single processors built from an entire silicon wafer, aim to keep that movement on the wafer instead, trading a cluster's modularity for on-chip bandwidth (section 2.2.1). Both designs answer the same constraint, the cost of moving data between chips, from opposite directions.

1.2.3 Communication dominance and the network wall

On a single GPU, a dense training step keeps its gradients in local memory, so matrix multiplications dominate step time and inter-node communication is zero. Once that workload spans racks in a warehouse-scale computer, every step must exchange gradients or activations across inter-server links before parameters can update, which can make communication the dominant cost and expose the network wall of the scale moment.

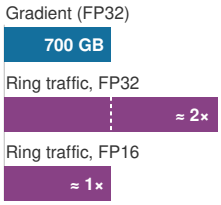

Ring AllReduce moves about twice the gradient per step; 16-bit gradients halve it.

Distributed data-parallel training synchronizes gradients across workers after each batch. For a model with 175B parameters, FP32 gradients occupy about 700 GB before any collective algorithm is applied. Ring AllReduce arranges the AllReduce exchange among neighbors in a logical ring. With N workers, each worker sends roughly $2(N-1)/N$ times the gradient tensor size and receives the same amount across the reduction and distribution phases (section 5.6.3 derives this factor). For large N the factor approaches two, so each worker moves about twice the FP32 gradient per step, and half that, about 700 GB, when mixed-precision training exchanges 16-bit gradients. Network traffic therefore depends on precision, worker count, collective implementation, and partitioning. On slower interconnects or partitions with little computation per communicated byte, it can consume a large fraction of each iteration rather than a negligible one. Notebook 1.1 later prices the gradient exchange of the 175B reference job in seconds per step.

↪ The α-β communication model, with the per-interconnect latency and bandwidth parameters that turn a link's specifications into a predicted transfer cost, is developed in section D.1.

Communication's share of the step grows as a fixed workload spreads across more devices, because each worker's local computation shrinks while its gradient exchange does not. Across illustrative tiers from 8 to 4,096 GPUs, compute can fall from 75 percent of the step to 20 percent as communication widens (figure 1.6). The fractions are assumed regimes rather than measurements. Gradient size, local batch size, partition topology, and overlap determine where the critical path actually lies, so it must be measured at each scale.

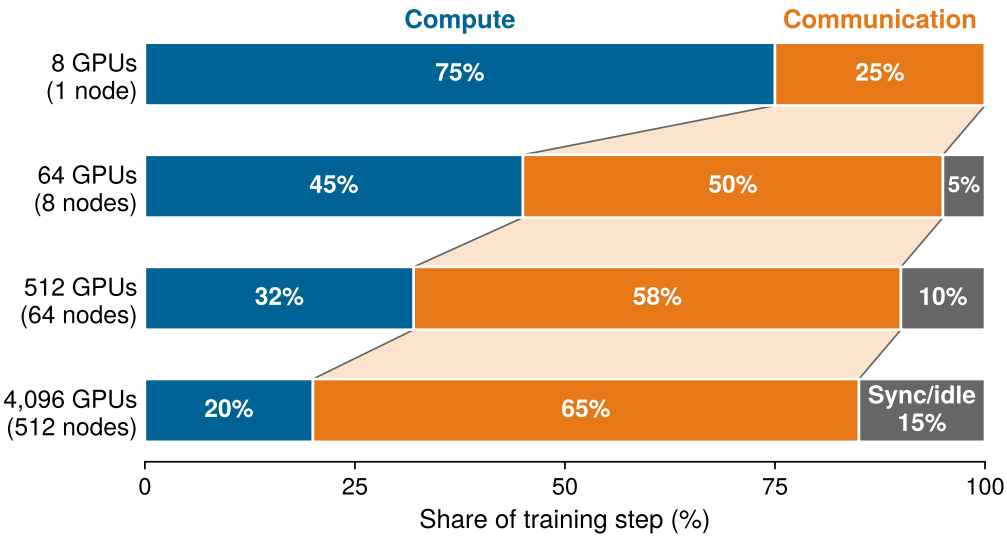

Figure 1.6: The Compute-Communication Timeline: Illustrative fractions show how exposed communication and synchronization can take a growing share of a training step as a fixed workload spreads across more devices. The fractions are assumed scenarios, not measurements or predictions for a particular model or cluster; they motivate measuring the critical path at each scale.

This exposed transfer cost explains why distributed training systems optimize communication so aggressively. Distributing work also distributes state that must move or be kept consistent. Horovod, for example, runs Ring AllReduce and fuses many small gradient tensors into fewer large transfers, so that the fixed cost of starting each exchange is paid fewer times (Sergeev and Balso 2018). Splitting each layer across devices, as Megatron-LM does with tensor parallelism (Shoeybi et al. 2019), and sharding optimizer state across workers, as ZeRO does (Rajbhandari et al. 2020), trade per-device memory for different traffic patterns. Gradient sparsification (sending fewer gradient bytes) and 3D parallelism (partitioning data, tensor, and pipeline dimensions across different fabric tiers) change how much data moves and where it moves. Chapter 6 develops these partitioning methods. Overlapping a transfer with independent computation instead reduces exposed waiting time without changing how many bytes move. The design decision at this stage is which cost prevents progress; selecting a particular mechanism requires the workload and fabric models developed later.

1.2.4 Communication intensity (the CI ratio)

If arithmetic intensity governs where a single accelerator stalls, **communication intensity (CI)** governs where a fleet stalls. On a single accelerator, the Roofline Model determines whether a kernel is compute-bound or memory-bound. At fleet scale, the same analysis extends to the network.

Equation 1.1 defines communication intensity as the ratio of data moved across the network to the operations performed locally. The numerator counts bytes and the denominator counts operations for the same worker and interval, such as one training step; the ratio therefore has units of bytes per operation.

$$\mathrm{CI} = \frac{M}{O} \tag{1.1}$$

where M is the number of bytes the worker moves across the network and O is the number of operations it performs locally in the same interval. A comparison must also state whether it counts sent bytes, received bytes, or both.

Low communication intensity describes compute-heavy workloads whose accelerators spend most of their time doing math, so scaling is comparatively easy. High communication intensity describes workloads whose bytes per operation approach or exceed the hardware's fabric balance, the ratio of per-worker network bandwidth to peak arithmetic throughput. Past that balance point, network transfers take longer than the local computation they accompany. Fabric balance plays the role of the ridge point with the ratio flipped, so a workload approaches the network limit as its communication intensity rises, whereas a kernel approaches the memory limit as its arithmetic intensity falls.

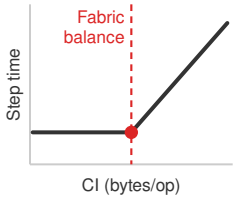

Past the fabric balance point, step time grows with communication intensity.

Splitting a fixed batch across still more accelerators shrinks the local operations O per worker while leaving gradient traffic M in place or adding collective overhead, which raises CI further and can make the step slower rather than faster. The balance point is configuration-specific because it depends on both local compute throughput and effective network bandwidth. When the workers' traffic must cross the fabric's narrowest cut, that effective bandwidth is each worker's share of the bisection bandwidth, so past the balance point the cut, rather than the accelerator, is the limit.

The 175B reference job gives the ratio a concrete value. Section 1.6.1.1 fixes that job as Archetype A's running instance, with 1,024 H100 GPUs, each data-parallel worker a replica of the model spread over 64 GPUs, and a global batch of 1,536 sequences per step. Each of the 16 replicas therefore processes 96 sequences, or 196,608 tokens, per step. Dividing both the replica's traffic and its work among its GPUs leaves the ratio unchanged, so it can be computed per GPU. Each GPU's share of the work is about six FLOPs per parameter per token, two in the forward pass and four in the backward pass (equation A.4), or 3.2×10^{15} FLOPs, while its 16-bit ring AllReduce of its gradient slice across the replicas moves about 10.3 GB. The communication intensity is therefore about 3.2×10^{-6} bytes per FLOP. An H100 on one 400 Gb/s InfiniBand link has a fabric balance of 5.1×10^{-5} bytes per FLOP, its link bandwidth divided by its peak arithmetic rate, so the job sits at about 6 percent of balance. At line rate and peak throughput, the gradient exchange takes only that fraction of the time the computation takes, which leaves the job far from the network wall and dominated by computation.

Spreading the same job over 8,192 GPUs moves it toward the wall. With the global batch fixed, each of the 128 replicas processes only 12 sequences per step, so each GPU's local work shrinks by

the same 8× as the cluster grows, to 4.0×10^{14} FLOPs. Its ring traffic barely changes, rising to 10.9 GB as the ring factor approaches two. Communication intensity climbs 8.5× to 2.7×10^{-5} bytes per FLOP, about 53 percent of the same balance, so the exchange now takes about half as long as the computation at peak and must be hidden by overlap or appear in every step. The cost of keeping replicas in step grows with the fleet because each accelerator has less local work to set against an exchange of nearly fixed size. These values are modeled rather than measured, and they count only the data-parallel exchange that crosses the fabric; tensor-parallel traffic stays on each node's NVLink. The ratio describes workload demand; the fabric and execution schedule determine how much of that demand delays progress. Methods that send fewer bytes, such as gradient sparsification, lower the ratio itself, whereas overlap leaves the ratio unchanged and only hides exposed time.

↪ The communication-computation ratio, and the threshold beyond which communication can no longer hide behind computation, are developed in section C.3.3.

Communication costs, together with the recovery costs quantified at the scale moment, describe how a fleet executes work. Neither explains what drives the relentless growth in fleet size, or how much work a model should receive. Both answers lie in empirical relationships between model quality and resource investment, relationships that have made warehouse-scale infrastructure an economic necessity rather than an engineering luxury.

1.3 AI Scaling Laws

The seven to eight orders of magnitude of training compute traced at the scale moment, from AlexNet to GPT-4-class estimates near 10^{25} FLOPs, were not an accident of ambition. Each order-of-magnitude increase in compute demanded a further expansion of the machine learning fleet, and organizations paid for those expansions because the return on them had become predictable. Predictable is not the same as free. Measured learning curves estimate the quality a training budget buys, while device throughput, execution efficiency, and the completion deadline determine how much fleet capacity that budget requires.

Rich Sutton's "bitter lesson" articulated the underlying principle: performance in machine learning is primarily driven by applying general methods at large scale rather than encoding human knowledge into algorithms (Sutton 2019). Scaling laws formalize this observation quantitatively. Model loss improves sublinearly as a power-law function of compute, dataset size, and parameters, $\mathcal{L}(X) \propto X^{-\alpha_{\text{scale}}}$ for a resource X (stated in full, with its irreducible floor, in theorem 1.1) (Kaplan et al. 2020; Hoffmann et al. 2022). Each equal-sized loss improvement therefore requires disproportionately more resources, a systems consequence this book later names as the empirical scaling law. Each scaling dimension (parameters, data, and compute) interacts with infrastructure constraints differently, making multi-dimensional efficiency optimization essential at production scale.

1.3.1 Empirical evidence for scaling laws

Over the longer history of AI, the growth rate of training compute broke sharply upward around 2010 (Sevilla et al. 2022). A fit to Epoch AI's records of notable models gives growth of about 1.5× per year before 2010, close to Moore's law, and of about 4.3× per year in the deep learning era (figure 1.7). The scale moment already decomposed that post-2012 climb into accelerator count, run duration, and delivered rate per accelerator (figure 1.1, panel b); that decomposition's lower rate of about 3.9× per year comes from fitting only the runs from 2012 through 2024 that report all three factors. When compute demand outpaces Moore's law by that much, misallocating a training budget becomes prohibitively expensive, turning empirical scaling laws from a research curiosity into an engineering necessity.

Kaplan et al. (2020) demonstrated that transformer-based language-model loss follows predictable power-law relationships with model parameters, dataset size (measured in tokens), and total computational budget. The result was not a license to increase every dimension blindly; it showed that model quality depends on coordinated allocation across those resources, a point later sharpened by Chinchilla's compute-optimal scaling analysis (Hoffmann et al. 2022).

Across parameter scales from 10^3 to 10^9, larger models are more sample-efficient, reaching any target loss with substantially fewer training tokens (figure 1.8, panel a). Each token costs a larger model more compute, yet an early-stopped larger model still attains the target loss with less total compute than a smaller model trained to its plateau (panel b). Compute-efficient training therefore stops far short of full convergence.

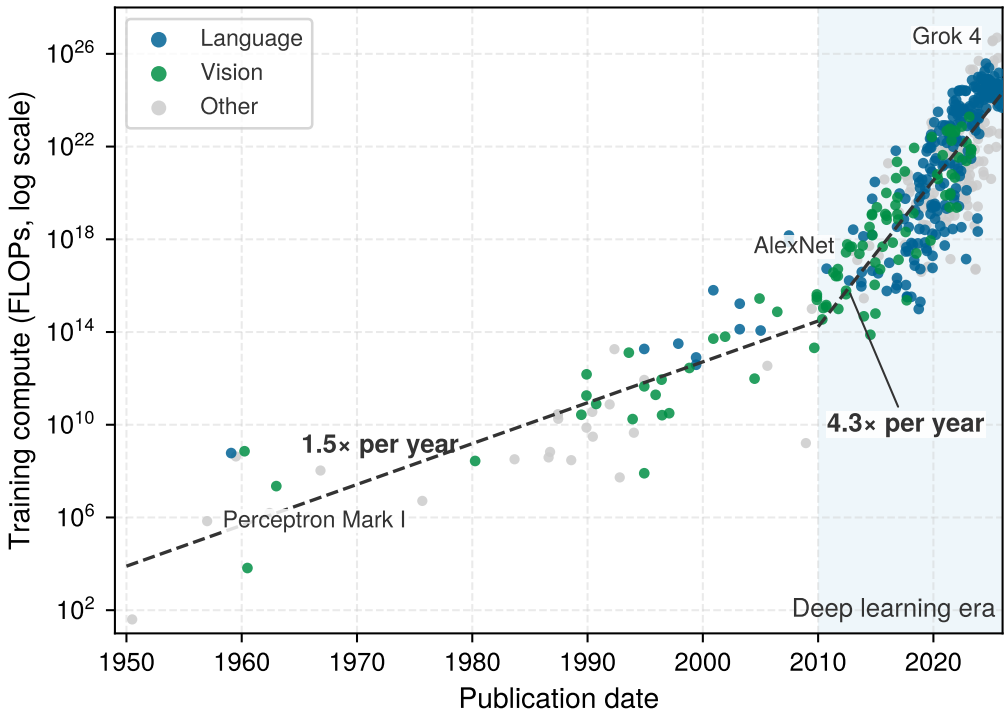

Figure 1.7: Model Training Compute Trends: Training compute of notable models published from 1950 through 2025, on a logarithmic axis, with one least-squares trend fitted to each era. Before 2010 the fitted rate is about 1.5× per year, close to Moore's law (a doubling about every 2 years); in the deep learning era it is about 4.3× per year, far steeper. The fastest stretch was steeper still. Within the 2012–2019 frontier subset measured by an earlier study (Amodei and Hernandez 2018), compute doubled approximately every 3.4 months. This trajectory explains why efficiency has become a strategic imperative rather than an optional optimization. Data from Epoch AI (2026c), licensed under CC BY 4.0; many values are Epoch AI estimates, and the era break at 2010 follows Sevilla et al. (2022).

Because each equal loss improvement costs disproportionately more, allocation grows more consequential as budgets grow. The question is how to divide a budget between the size of the model and the training data it consumes.

1.3.2 Compute-optimal resource allocation

A team that starts from a balanced allocation and spends the same training-compute budget on a model twice as large, with half the training tokens,[5] ends up with a higher-loss model that also costs twice as much memory and bandwidth to serve. Empirical studies of large language models (LLMs) show that, for a fixed computational budget and within the studied scaling regime, balancing model size and dataset size minimizes training loss.

Holding the budget fixed and varying model size, so that a larger model receives fewer tokens, traces a U-shaped IsoFLOP curve of training loss against parameter count, whose valley marks the best model size for that FLOP budget (figure 1.9, panel a). As the budget grows, the valley moves toward larger models, and the optimal parameter count and the optimal number of training tokens rise along coordinated power laws (panel b), so compute-optimal training advances model size and token volume in tandem. Projected to the budget of Gopher, an earlier DeepMind language model, the fits call for 63B parameters trained on 1.4T tokens, about 5.3×10^{23} FLOPs in total,[6] which is how compute-optimal frontiers guide resource allocation before cluster execution begins.

Within the studied regime, the Chinchilla result indicates that model size and training tokens should scale approximately proportionally for compute-optimal training, with a ratio on the order of 20 training tokens per parameter (Hoffmann et al. 2022). The allocation is a trade along a fixed product. Training a dense transformer costs about $O_{\text{train}} \approx 6PD$ FLOPs for P parameters and D

[5] **Tokens**: Subword units produced by algorithms like byte-pair encoding, which iteratively merges the most frequent character pairs in a corpus. Token vocabulary size creates a direct systems trade-off: larger vocabularies can reduce sequence length for some languages and domains, but they expand the embedding table roughly in proportion to vocabulary size. GPT-3 used a 50,257-token vocabulary; moving to 100K+ tokens in newer models therefore trades tokenization efficiency against additional embedding memory.

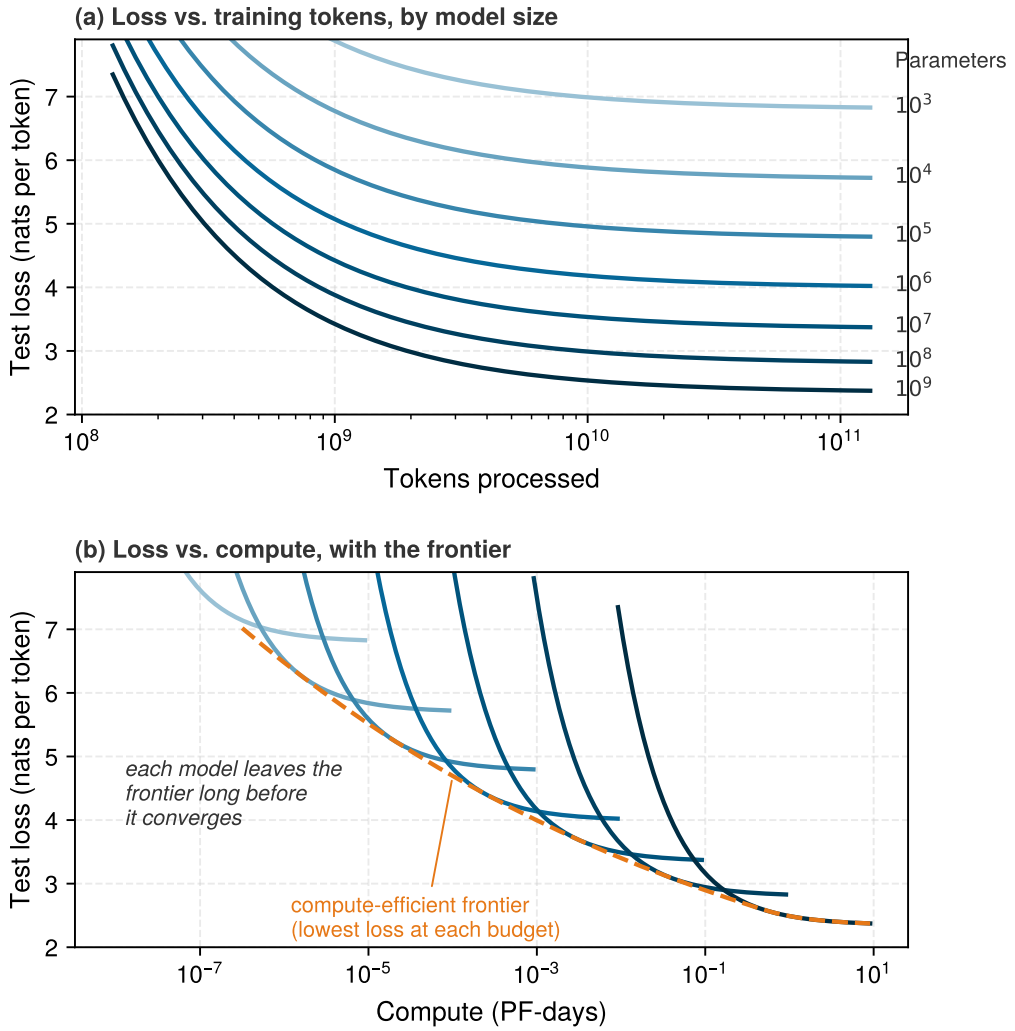

Figure 1.8: Scaling Laws and Compute Optimality: Test loss predicted by the fitted law of Kaplan et al. for models from 10^3 to 10^9 parameters. (a) Larger models reach any given loss with fewer training tokens and settle at a lower plateau. (b) The same curves against training compute; the dashed line is their lower envelope, the lowest loss any model size reaches at each budget. Each model leaves the envelope long before it converges, so compute-efficient training stops far short of convergence. The curves are drawn from the published fit of loss against parameter count and training steps at the paper's default batch size, not from training logs, and begin where the fit's early-training term falls to 5 nats. After Kaplan et al. (2020).

training tokens, two per parameter per token in the forward pass and four in the backward pass (equation A.4). A fixed budget therefore fixes the product of model size and tokens, and the ratio chooses the point along it. The projection in figure 1.9, 1.4T tokens for 63B parameters, or about 22 tokens per parameter, sits close to that ratio. This corrected earlier guidance that favored making models larger while training them on comparatively fewer tokens. The objective is loss at a fixed training-compute budget. A deployment that serves the model repeatedly must also account for inference cost; its preferred allocation need not minimize training compute alone.

When sizing GPT-3, the authors used scaling-law extrapolations to choose model size and training data (Brown et al. 2020; Kaplan et al. 2020). They trained an established transformer architecture with 175B parameters and approximately 300 billion tokens, about 1.7 tokens per parameter. Against the roughly 20 tokens per parameter of the later Chinchilla results, that point was undertrained relative to compute-optimal token allocation (Hoffmann et al. 2022).

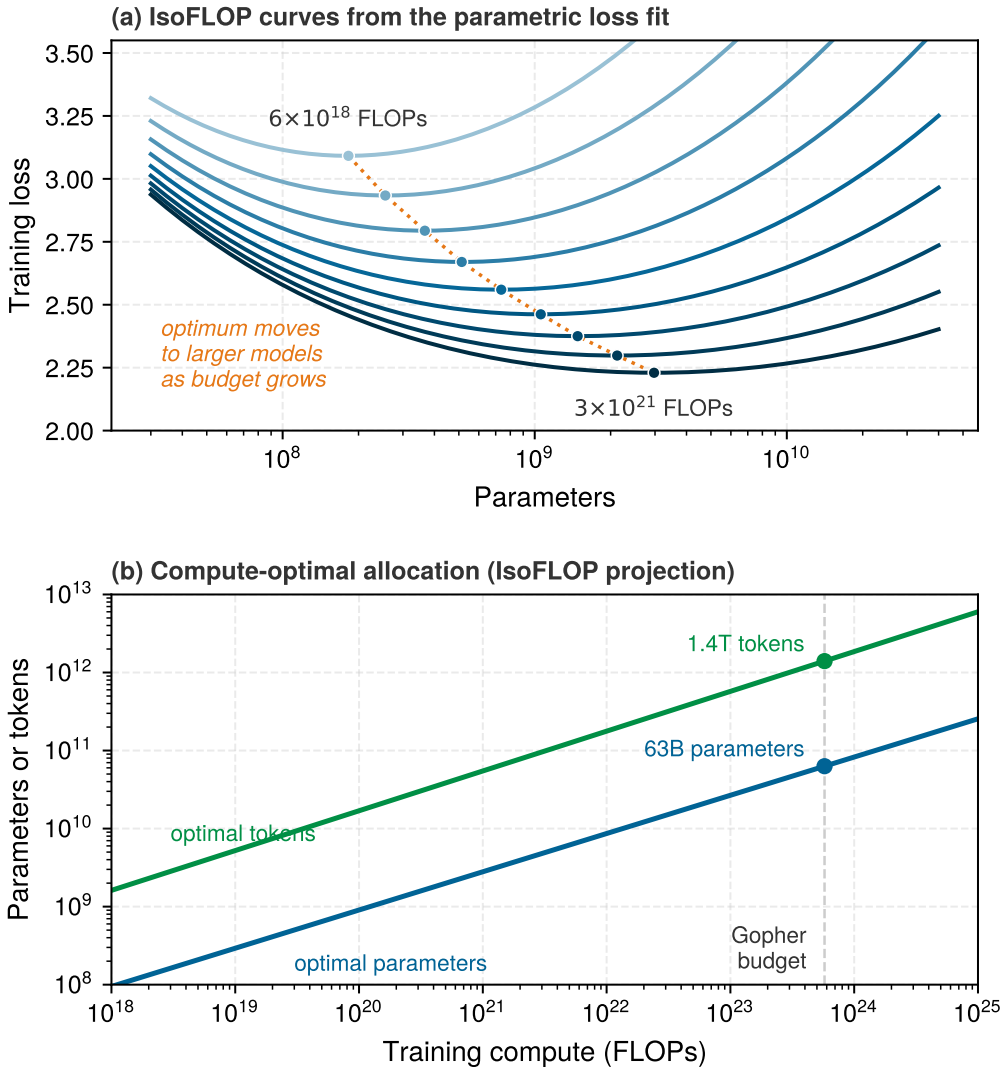

Figure 1.9: Optimal Compute Allocation: For a fixed computational budget, language-model loss depends on how the budget is divided between model size and training data. (a) IsoFLOP curves for nine budgets from 6×10^{18} to 3×10^{21} FLOPs, drawn from the paper's parametric loss fit; each valley marks the best model size for its budget, and the valleys move toward larger models as the budget grows. (b) Compute-optimal parameter and token counts from the paper's IsoFLOP analysis, projected to the Gopher budget, where they reach 63B parameters and 1.4T tokens. The two panels use different fitting methods from the same paper, so their optima do not coincide exactly. After Hoffmann et al. (2022).

The practical decision is how to divide a fixed budget among model size, data, and training duration, and a compute-optimal allocation answers it only in useful computational work. Converting that work into elapsed time requires a separate systems model. Communication can grow relative to per-worker computation as a fixed workload is split more finely, creating bandwidth bottlenecks that reduce effective utilization. At multi-thousand-GPU scale, utilization depends on partitioning and interconnect. Narayanan et al. (2021) report Megatron-LM runs on 3,072 A100 GPUs at 52 percent of peak device throughput and note that slower interconnects or communication-heavy partitions hinder scaling.

1.3.3 Power-law scaling and operational regimes

Even when a cluster sustains high utilization on a compute-optimal run, scaling parameters and tokens in lockstep works only while both resources remain available and the loss curve stays well above its irreducible floor. That floor, left out of the power law as stated with the bitter lesson, explains why the same fleet can enter different operational regimes. The following analysis takes three cuts through these regimes, each with its own fleet consequence. The first cut is how compute and data are allocated, which decides the resource to buy next. The second is where diminishing returns set in with dataset size, which decides when storage and data pipelines matter more than accelerators. The last is when in the model's lifecycle computation is spent, which decides the part of the fleet, training or serving, that carries the load.

Theorem 1.1: Power-law scaling formulation

Scaling laws are expressed formally as power-law relationships. The general formulation is:

$$\mathcal{L}(X) = A_{\mathcal{L}} X^{-\alpha_{\text{scale}}} + C_{\mathcal{L}}$$

where loss $\mathcal{L}$ decreases as resource quantity X increases, following a power-law decay with rate α_{scale}, plus a baseline constant $C_{\mathcal{L}}$, the irreducible loss (a constant of the fit, not one of the C^3 labels). Here, $\mathcal{L}(X)$ represents the loss achieved with a resource quantity such as parameters, tokens, or compute; $A_{\mathcal{L}}$ and $C_{\mathcal{L}}$ are task-dependent constants; and α_{scale} is the scaling exponent that characterizes the rate of performance improvement. A larger value of α_{scale} signifies more efficient performance improvements with respect to scaling.

For a fleet, the exponent is the price list. Halving the reducible term $A_{\mathcal{L}} X^{-\alpha_{\text{scale}}}$ requires multiplying the resource X by $2^{1/\alpha_{\text{scale}}}$, so the smaller the exponent, the larger the multiple of accelerators, tokens, or training time that each further step in quality demands. These power-law predictions find empirical validation in the family of learning curves plotted in figure 1.10. Across six model scales spanning from 393.2K to 708M parameters, test loss descends predictably as dataset size grows from 10^7 toward 3×10^{10} tokens, with smaller architectures flattening into parameter-capacity plateaus while larger models sustain power-law error reduction.

1.3.3.1 Resource-constrained scaling regimes

The first cut follows which resource runs short. The capacity plateaus show why collecting more tokens for a project whose cluster cannot train a larger model strands data in storage. The opposite imbalance, buying more accelerators for a project whose dataset is already exhausted, instead risks wasting hardware on overfitting. Applying scaling laws in practice is therefore a scarcity diagnosis that asks which resource prevents balanced growth first. Compute budget, data availability, and model size create three regimes with different optimal responses.

In a compute-limited regime, compute scarcity restricts scaling potential despite abundant training data; measured learning curves guide the allocation between model size and training tokens within that budget. In a data-limited regime, computational resources exceed what the available dataset can support, meaning that increasing model size is useful only if validation measurements show that the available data can support additional capacity. Between these extremes lies the optimal regime, where compute and data grow in balance following compute-optimal scaling laws, as DeepMind demonstrated with Chinchilla by outperforming much larger models through proportional scaling of model size and training data (Hoffmann et al. 2022). Recognizing which regime governs a given project prevents common inefficiencies: over-parameterized models with insufficient training data, or under-parameterized models that waste available compute. For the fleet, the diagnosis decides where the next investment goes. A compute-limited project needs accelerator-hours, while a data-limited one gains more from the storage and data pipelines of chapter 4 than from additional accelerators.

The second cut through the operational regimes follows the return on additional data as a dataset grows, which decides when buying data stops paying. On log-log coordinates, dataset scaling passes through three regions (figure 1.11). In the initial Small Data Region, generalization error

6 **FLOPs vs. FLOP/s**: FLOPs denotes total computational work (operations performed), while FLOP/s denotes hardware throughput (operations per second); confusing the two leads to incorrect cost estimates. GPT-3 required 3.14×10^{23} FLOPs of work; an A100 delivers 312 TFLOP/s of throughput. Dividing work by peak throughput gives ideal wall-clock time, the roughly 32 years on one A100 computed at the scale moment. Actual time is approximately work$/(\eta_{\text{hw}} R_{\text{peak}})$, where the hardware utilization η_{hw} is the achieved fraction of peak throughput, so utilization below one multiplies ideal time by $1/\eta_{\text{hw}}$.

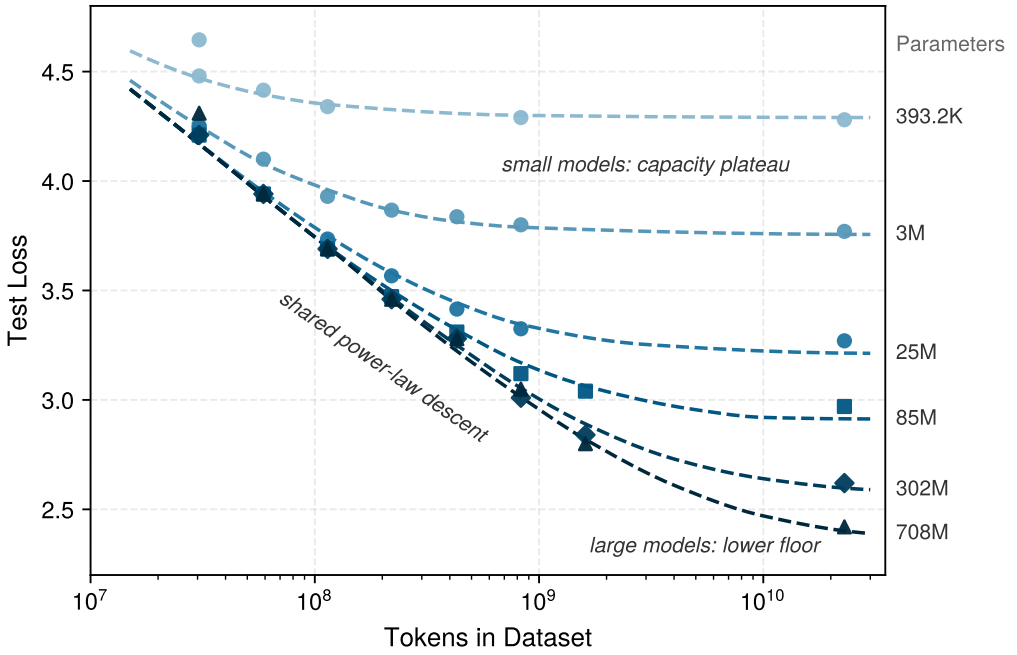

Figure 1.10: Loss vs. Dataset Size across Model Scales: Test loss curves show how models of different sizes (393.2K to 708M parameters) benefit from increased training data. Larger models achieve lower loss, but all curves exhibit diminishing returns at high token counts. Points and curves are digitized from Figure 4 (left) of Kaplan et al. (2020), so values carry the precision of a reading from the published plot.

lingers near the upper Best Guess Error ceiling, the error of a model that has learned almost nothing from its data. As data volume expands into the central Power-Law Region, error drops sharply along a constant slope where data investment yields predictable returns. Finally, in the Irreducible Error Region, the curve flattens asymptotically against the Bayes noise floor, where further dataset expansion provides negligible gain. That floor is the error that remains even for a perfect predictor, because the inputs do not fully determine the labels. The systems consequence depends on the region. In the Power-Law Region, finite high-quality data, not compute, can become the binding constraint, which is why data efficiency complements brute-force scaling. In the Irreducible Error Region, more data of the same kind stops paying.

An observed plateau does not by itself establish that the irreducible floor has been reached, however. Limited coverage, insufficient model capacity, and incomplete optimization can also flatten a learning curve, so the engineering task is to rule out these remediable limits before purchasing more data or compute.

↪ The data wall, and the storage traffic a fleet incurs when it generates, filters, and versions synthetic training data, are developed in section 4.7.2.

The third cut follows lifecycle stage, because additional compute can be spent on the foundation run, on adapting the model afterward, or on each request in deployment (figure 1.12). Pretraining scaling (orange curve) increases the resource investment in the foundation run. Posttraining scaling (blue curve) adapts an existing model through fine-tuning or preference optimization (training on human rankings of its outputs). Test-time scaling (green curve, "long thinking") spends additional computation on a particular request through intermediate generation (writing out reasoning steps before the answer), search, or sampling, trading serving latency and throughput for a possible quality gain. Each phase stresses a different part of the fleet. Pretraining is one long, failure-sensitive foundation run whose scaling-law gains must be realized in wall-clock time. Posttraining replaces that single large job with many smaller, coordination-heavy runs. Test-time scaling moves computation into the serving path, where it competes with latency targets. The schematic curves do not establish a universal efficiency ranking among these choices, so the full treatments belong with their workloads. Test-time compute at serving is treated in section 10.6.9. A resource plan

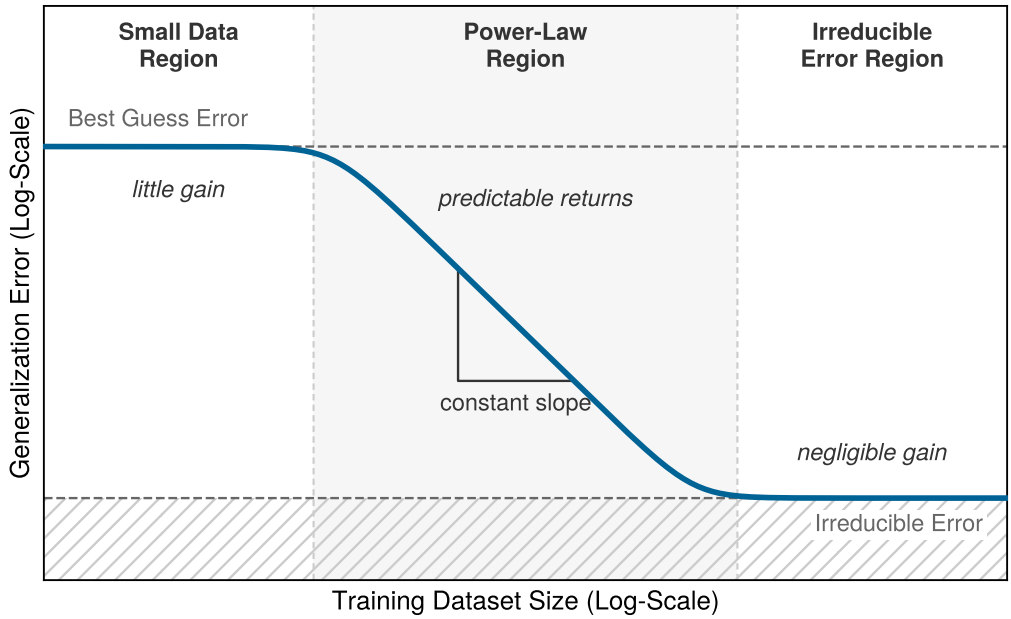

Figure 1.11: Data Scaling Regimes: Generalization error as a function of dataset size exhibits three distinct phases. In the small-data regime, error stays near the best-guess ceiling; the power-law regime demonstrates predictable error reduction; the irreducible-error regime represents the fundamental noise floor. Adapted from (Hestness et al. 2017); the axes carry no measured values.

must compare measured task improvements and costs before choosing where to spend the next increment of compute.

↪ The multi-model coordination that posttraining alignment runs require is developed in section 6.7.1.

1.3.4 Sustainability and cost implications

GPT-3's training run alone occupied about 148,000 GPU-days and drew an estimated 1,287 MWh of electricity (Patterson et al. 2021), and later frontier runs use more accelerators for longer. While scaling laws trace a path to lower loss, they also expose rapidly escalating resource demands whose training and deployment costs grow faster than the quality gains they buy.

Where that electricity is generated matters as much as how much is drawn. The same training run can differ by roughly 41× in CO_2 emissions depending on grid location, from about 20 g/kWh of CO_2 for Quebec hydropower to about 820 g/kWh for Poland coal. The siting of the cluster is therefore a first-order infrastructure decision. Financial cost scales with the same accelerator-days and megawatt-hours, which limits who can do large-scale research at all.

↪ How site selection and grid carbon intensity set the emissions of the same training run is developed in section 15.2.

Scaling laws do not guarantee unbounded improvement. Each incremental performance gain must be weighed against its resource cost. As systems approach practical scaling limits, the emphasis shifts from scaling alone to efficient scaling: balancing performance, cost, energy consumption, and environmental impact.[7]

[7] **Scaling law saturation**: Power-law relationships with a baseline loss term approach that floor rather than predict unbounded loss reduction. Downstream task utility may also stop improving even while the fitted loss decreases. That gap between loss and utility motivates testing data efficiency or model compression against task utility and resource cost; neither technique guarantees that a genuine information limit can be overcome.

1.3.5 When scaling laws break down

Doubling a model's parameter count without expanding its training dataset or interconnect capacity can stall loss improvements even as cluster cost doubles, because power-law fits hold only within the balanced regime in which they were measured. Most breakdowns at the edge of that regime share one root cause, growth in one dimension that outpaces the others. Hoffmann et al. (2022) show that many large language models were undertrained because model size grew faster than training-token budgets, and that compute-optimal training requires scaling parameters and tokens together. Hardware ceilings belong to the same family, because memory bandwidth, interconnect capacity, and I/O throughput cap how far a larger model can be distributed at all.

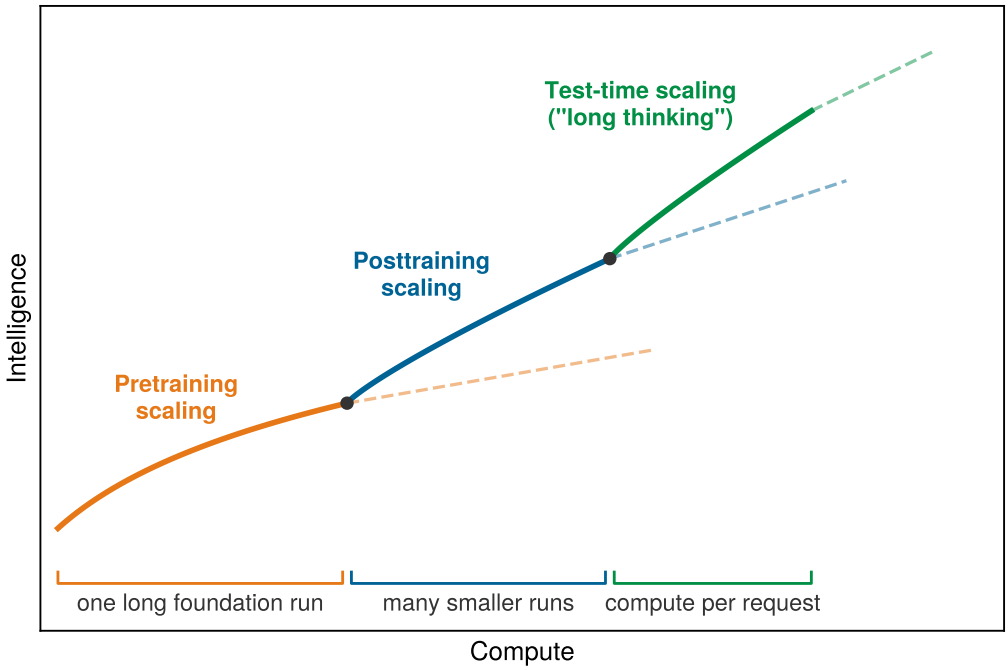

Figure 1.12: The Three Scaling Phases: A schematic of three points in the lifecycle where additional computation can raise capability. Pretraining scales with compute and data (Kaplan et al. 2020). Posttraining (reinforcement learning from human feedback, instruction tuning) refines behavior (Ouyang et al. 2022). Test-time scaling (chain-of-thought, search) spends additional computation on each request at inference (Snell et al. 2025). The curves are illustrative and carry no measured values.

Each imbalance requires different evidence. Training measurements can reveal an unsuitable optimization schedule. Held-out evaluation can distinguish improved generalization from memorization. Execution traces can reveal whether bandwidth or I/O prevents the fleet from completing the planned work. A lower benchmark loss also does not establish robustness to adversarial or out-of-distribution inputs. Each possible breakdown type pairs with a candidate cause and a representative scenario (table 1.2). The pairing guides investigation; it does not infer a cause from a plateau alone.

Table 1.2: Possible Scaling Breakdowns: The rows describe risks and representative scenarios, not inevitable consequences of increasing one resource. Training and validation measurements establish which failure mode, if any, applies to a particular allocation.

Dimension scaled	Type of breakdown	Underlying cause	Example scenario
Model Size	Overfitting	Model capacity exceeds available data	Billion-parameter model on limited dataset
Training Data Size	Diminishing Returns	Saturation of new or diverse information	Scaling web text beyond useful threshold
Compute Budget	Underutilized Resources	Insufficient training steps or inefficient use	Large model with truncated training duration
Imbalanced Scaling	Inefficiency	Uncoordinated increase in model/data/compute	Doubling model size without more data or time
All Dimensions	Saturation	Exhaustion of learnable patterns in the domain	No further gains despite scaling all inputs

Checkpoint 1.2: Applying scaling laws

These questions check whether scaling laws are being used as resource-allocation tools:

- ☐ **Compute vs. data regime**: identify the learning-curve evidence needed to choose between a larger model trained for fewer steps and a smaller model trained for more steps under fixed compute and limited data.
- ☐ **Scaling breakdown**: explain the risks of doubling model parameters at fixed dataset size, and what training and validation evidence would distinguish undertraining from overfitting.
- ☐ **Chinchilla optimality**: explain why Chinchilla optimality serves as a resource-allocation target and identify the minimized quantity.
- ☐ **Compute vs. capability**: classify the claim as true or false: scaling laws predict that doubling compute always yields a doubling of model capability.

1.3.6 Combining efficiency with scaling

Data saturation, infrastructure bottlenecks, and diminishing returns set hard limits on what brute-force scaling can achieve. The same constraints, however, point toward solutions along three interconnected dimensions. When compute is the scarce resource, algorithmic efficiency uses techniques such as sparsity (computing with only a fraction of the weights) and distillation (training a small model to imitate a large one) to extract more capability per FLOP. When high-quality tokens run short, data efficiency uses curriculum learning (ordering examples by difficulty) and active selection (choosing the most informative examples) to extract more information per training example. When the cluster itself stalls on transfers or barriers, systems efficiency uses communication hiding and pipeline overlap, which run transfers behind useful computation, to extract more utilization per accelerator. Because each dimension addresses a different breakdown condition from table 1.2, identifying the limiting resource determines whether efficiency work should target the algorithm, the data, or the execution system. Systems efficiency is the subject of chapter 5 and chapter 6; the latter also develops algorithmic levers such as reduced precision and sparse mixture-of-experts execution (running each token through only a few of a model's subnetworks). Data efficiency under device budgets returns in section 11.5.

Scaling laws and efficiency techniques determine how much computation large models require and how well the fleet uses it. They do not, however, address the physical and logical constraints that the fleet itself imposes. Hardware fails, bandwidth saturates, distributed systems face fundamental impossibility results, and societal impact at scale demands governance. These constraints cannot be optimized away; they must be engineered around.

1.4 Constraints of Scale

Scaling laws set how much work a frontier run must perform, but a cluster sized for that budget can still stall every few hours on a failed node, halt at a synchronization barrier during a network partition, or be blocked from release because its training pipeline lacks an audit trail. Those failure modes extend beyond the three walls of the scale moment into three broader categories of irreducible constraint: physical (the memory, network, and energy walls themselves, plus the reliability gap that opens as fleet size grows), logical (distributed systems face impossibility results like the CAP theorem), and societal (scale amplifies the impact of every technical decision). For an initial C^3 diagnosis, the distinctions matter because a slow step may reflect a communication limit or coordination loss, while an otherwise fast design may exceed its energy or governance budget.

1.4.1 The reliability gap

A server node that is unavailable for less than an hour a year seems dependable in isolation, yet an unreplicated synchronous job spanning 25,000 GPUs can advance only while every required device is healthy at the same instant. This divergence opens the **reliability gap**: high component reliability in isolation does not ensure that every component required by an unreplicated job will be available

simultaneously. The scale moment measured the gap as a rate, how often the job is interrupted; availability measures it as a probability, how likely it is that every required component is up at a given instant. Both views rest on the same independence assumption and degrade the same way as the job grows. For that 25,000-GPU job, equation 1.2 multiplies the component availabilities under an independence assumption to give the probability that all required components are available at the same instant (A_{all}):

$$A_{\text{all}} = (A_{\text{component}})^N \tag{1.2}$$

Here, N is the number of independent components that the job requires simultaneously, and $A_{\text{component}}$ is the availability of one component. The exponent carries the engineering implication. Every added required component lowers the probability that the full set is available, so reliability engineering becomes a scaling problem rather than a cleanup task.

For a single component, 99.99 percent availability corresponds to 52.56 minutes of unavailability per 365-day year. This is a time fraction, not a count of failures. If each node has an illustrative 99.9 percent availability and an unreplicated job requires every node, the probability that all 1,000 required nodes are available simultaneously is only 36.8 percent. Scale that job to 10,000 required nodes, and this probability drops to 0.0045 percent. Adding a nine to every node only postpones the collapse, because the exponent N eventually overwhelms any per-node availability below one (figure 1.13).

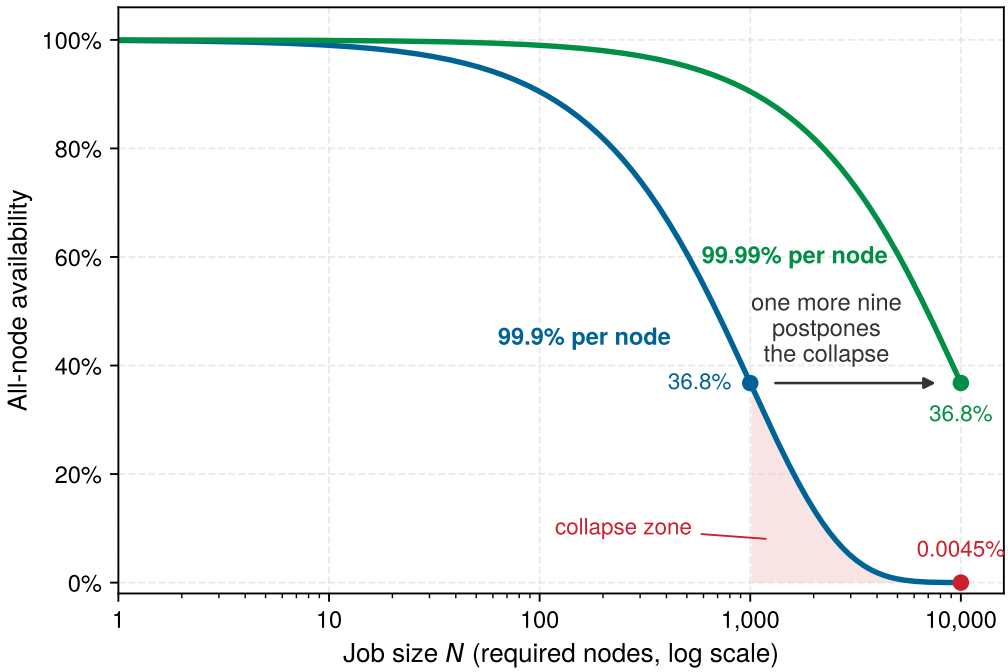

Figure 1.13: The Reliability Collapse: The simultaneous availability of every node required by an unreplicated job, shown as a function of job size for two per-node availability levels under an independence assumption. At 99.9 percent per-node availability (blue), all-node availability is only 36.8 percent for a 1,000-node job and approaches zero for a 10,000-node job. Even 99.99 percent per-node availability (green) yields 36.8 percent at 10,000-node scale.

Because simultaneous all-node availability drops toward zero as job size expands into the thousands (the collapse zone of figure 1.13), an unreplicated synchronous workload cannot rely on unbroken component uptime. Section E.1 replaces this single per-node availability figure with the full failure-probability model, from component failure rates to the chance that a job fails before it finishes. The engineering lesson is that failure is the common case. At scale, resilience replaces prevention as the governing objective, and the system trades absolute uptime for recovery speed, checkpoint quality, and automated repair.

↪ The shift from keeping every component running to a fleet that recovers automatically when components fail is developed in chapter 7.

Checkpoint 1.3: The scale mandate

These questions check whether the scale mindset has shifted from local execution to distributed constraints:

- ☐ **Communication intensity**: explain in CI-ratio terms why adding more nodes (N_{nodes}) to a high-CI workload can slow training rather than speed it up.
- ☐ **Reliability gap**: explain the reliability gap that prevents a 10,000-GPU cluster from remaining "perfectly healthy."
- ☐ **Arithmetic vs. communication**: distinguish arithmetic intensity (local memory) from communication intensity (global network).

1.4.2 Physical boundaries of distributed state

Scale forces distribution, because no single machine provides the compute required for the largest models and no centralized system can collect every data source that global user bases generate. Coordinating computation across physically separated machines connected by finite-bandwidth networks, however, creates logical constraints that no engineering can eliminate. The first appears inside data centers, where synchronous training must choose how much consistency to preserve under partitions and failures. The second appears at the edge, where the same choice returns under harsher conditions because the devices themselves are unreliable, power-limited, and policy-constrained.

1.4.2.1 The CAP theorem reality

When a worker becomes unreachable in the middle of a distributed training step, the runtime faces an immediate choice: pause the entire job until the worker returns or is replaced, or keep updating the model without its contribution. The **CAP theorem**[8] states the same kind of forced choice for replicated data. It establishes that a distributed system on a network with no timing bounds, once a partition splits that network, cannot guarantee both atomic **consistency** (every read returns the latest completed write) and **availability** (in the theorem's sense, every request to a running node receives a response). **Partition tolerance** describes the network condition under which that trade-off applies, not a third interchangeable property that a system can choose to omit.

[8] **CAP (consistency, availability, partition tolerance) theorem** (Brewer's theorem): Brewer formulated the conjecture (Brewer 2000), and Gilbert and Lynch (2002) later proved it for asynchronous networks, meaning networks with no timing bounds.

Training is not a replicated database, and gradient freshness is not the theorem's atomic consistency, so the theorem does not apply directly, yet training faces the same forced choice under a fault. Synchronous training chooses a shared current model state, so the entire step halts when a required worker becomes unreachable. Asynchronous training makes the opposite trade, continuing through stragglers or unavailable workers at the cost of applying updates computed on stale model versions. Federated learning pushes that availability choice to the edge by letting intermittent devices train local models that diverge temporarily between rounds.

1.4.2.2 Edge distribution constraints

The coordination challenges discussed so far assume data-center distribution, where machines run in managed facilities. At the network edge, these challenges intensify along every dimension. Billions of smartphones and IoT devices operate in uncontrolled environments with unreliable connectivity and limited power. Privacy, consent, or policy constraints may prevent raw data from leaving the device, making federated learning a natural architecture (McMahan et al. 2017). Intermittent connectivity requires a participation policy, which may select currently available devices for a round or admit delayed updates under stated rules. Keeping raw data local does not by itself protect information exposed through model updates. Device budgets and the threat model determine whether additional compression, privacy mechanisms, or local execution are appropriate.

1.4.3 Governance and systemic risk at scale

Scale and distribution amplify impact beyond engineering. When a system serves billions of users, a technical bug becomes a societal risk. This amplification increases the reach of governance failures; smaller deployments still require controls appropriate to their uses and potential harms. Governance

functions as the **control plane** of the machine learning fleet, the layer that decides what the fleet may optimize and deploy, not as a set of external rules.

That control plane first has to defend the fleet against security threats that grow more rewarding to attackers at production scale. Model extraction attacks steal proprietary intellectual property through normal API queries, whereas data poisoning corrupts the training pipeline from the inside to alter learned behavior; one objective is a backdoor, a hidden behavior that activates on a chosen input. At fleet scale, these threats become economically attractive targets for sophisticated attackers. Defending the fleet requires systematic approaches: access controls, differential privacy (calibrated noise that bounds how much any single training record can influence what the model reveals), and continuous behavioral monitoring that go far beyond traditional perimeter security.

↪ Defenses against data poisoning, model extraction, and gradient inversion across distributed fabrics are developed in chapter 13.

The same control plane must make the fleet auditable. Systems operating at scale attract regulatory attention. Regimes such as the EU AI Act and local privacy laws make **auditability**, the ability to reconstruct after the fact which data, code, and decisions produced a model's behavior, an architectural requirement for many deployments. Auditing whether a model trained on 10,000 GPUs ingested prohibited data requires end-to-end data lineage tracking, supported by controls that capture ingestion and subsequent transformations. A lineage record cannot establish the absence of inputs that were never recorded. Generating a human-interpretable explanation for a recommendation served within a tail-latency budget of tens of milliseconds, a limit on how long the slowest requests may take, demands techniques designed into the serving architecture. Meeting these requirements calls for technical capabilities (audit trails, bias testing, and consent management) built into the infrastructure from the outset.

Beyond legal compliance, the machine learning fleet carries ethical obligations. Recommendation algorithms shape public discourse; hiring algorithms affect livelihoods. These systems rarely fail loudly; harm accumulates silently through bias and polarization, without a crash log. Bias in a widely reused foundation model can spread to the many systems built on it. Responsible AI turns fairness, transparency, and accountability requirements into operational invariants with defined evidence, owners, and responses. A violation may require withholding a release, routing affected decisions to review, or suspending an unsafe operation; the response must contain the relevant harm.

↪ The assurance lifecycle that turns responsible-AI requirements into evidence, accountable owners, and the authority to launch or stop a system is developed in section 16.2.

Together, these physical, logical, and societal constraints define the operating envelope of the fleet. Turning them into actionable systems decisions requires a diagnostic taxonomy that can isolate whether execution time is lost to local arithmetic, wire transit, or orchestration overhead.

1.5 The C^3 Taxonomy and Foundations of Scale

An idle accelerator does not identify its own bottleneck. It may be waiting for local operands, a gradient transfer, or a failed peer. The C^3 lens introduced at the scale moment separates these explanations and gives the preceding examples a common diagnostic language. On a single machine, diagnosing performance follows the **Data · Algorithm · Machine (D·A·M) taxonomy**. In that framework, Data governs what information the model learns from; the Algorithm defines the mathematical operations to execute; and the Machine supplies the silicon and memory hierarchy executing those operations.

When workloads stretch across thousands of interconnected nodes, the execution bottleneck shifts. Data-center hardware is procured and cabled around three physical pillars: Compute, Network, and Storage. At runtime, however, distributed execution spends wall-clock time in three distinct operational dimensions: Compute (C_1), Communication (C_2), and Coordination (C_3). Together, they form the C^3 taxonomy. Storage is not absent from this runtime model. It serves as the persistent physical substrate spanning all three dimensions: an extension of the node memory hierarchy in Compute (C_1), a high-throughput streaming pipeline across Communication (C_2), and the durable checkpointing engine required for fault-tolerant Coordination (C_3).

↪ The D·A·M taxonomy as a complete single-node framework, with its diagnostic summary, roofline bound, and scorecard, is developed in Appendix A.

Each dimension leaves different evidence. Suppose a training job's step time grows by half after its cluster doubles. If time inside each node's kernels grew while their model FLOPs utilization (MFU), the fraction of peak arithmetic throughput spent on the model's own operations, fell and memory-bandwidth counters saturated, local execution is responsible. If link utilization and collective timing rose, transfers are; if barrier waiting and recovery records grew, coordination is. The three diagnoses call for different interventions, and optimizing the wrong mechanism leaves the end-to-

end bottleneck in place. As in single-machine systems, the governing limit is the *binding constraint*: the bottleneck whose relaxation yields an end-to-end reduction in training time or serving latency. The C^3 taxonomy is therefore an operational diagnostic loop: identify the binding axis, apply the appropriate systems intervention, and remeasure, because relaxing one limit frequently shifts the bottleneck to another.

Definition 1.2: The C^3 taxonomy

C^3 taxonomy is a diagnostic framework that classifies distributed machine learning execution time and system bottlenecks along three operational dimensions: Compute (C_1, the math), Communication (C_2, the wire), and Coordination (C_3, the logic). The goal is to identify which dimension binds fleet-scale performance.

1. **Significance**: The diagnostic power is concrete even before inspecting complex distributed traces. If accelerators are computing at low MFU because tensor cores starve for local operands or un-fused kernels, the binding axis is Compute (C_1). If accelerators sit idle waiting for AllReduce gradient transfers or activation tensors across the network fabric, the binding axis is Communication (C_2). If progress halts due to straggler-induced barrier waiting, node recovery rollbacks, or preemptions under gang scheduling (which allocates all of a job's workers at once or none of them), the binding axis is Coordination (C_3).
2. **Distinction**: Unlike data center procurement categories (Compute, Network, Storage), which classify physical hardware assets, the C^3 taxonomy classifies runtime execution cost. Storage does not execute model arithmetic; rather, it serves as the persistent physical substrate across all three dimensions (local NVMe in C_1, remote dataset streaming in C_2, and durable checkpointing in C_3).
3. **Common pitfall**: Treating the three dimensions as independent variables. Upgrading to faster accelerators (C_1) often converts a compute-bound job into a communication-bound job (C_2). Sharding model states to eliminate memory redundancy (C_1) increases collective communication volume (C_2) and state coordination complexity (C_3). The dimensions move together.

1.5.1 The three dimensions of the fleet

Each dimension of the C^3 taxonomy governs a distinct physical layer of the fleet and exposes a different operational bottleneck.

1.5.1.1 Node compute substrate

When operands cannot reach a node's arithmetic units fast enough to keep tensor cores busy, the memory wall binds Compute (C_1, the math). This axis covers execution within the server node across arithmetic units, high-bandwidth memory, and local Non-Volatile Memory Express (NVMe) storage. Improving MFU on this axis requires reducing local data movement, for example through operator fusion, rather than merely increasing peak arithmetic throughput.

↪ The memory wall, from high-bandwidth memory (HBM) bandwidth to the Roofline Model that locates it for one accelerator, is developed in section 2.3.

1.5.1.2 Network communication backbone

When gradient, activation, or dataset-shard transfers saturate a cluster's finite bisection bandwidth, they reach the network wall and expose communication time (T_{comm}) during which accelerators wait for remote data. Communication (C_2, the wire) captures this data movement between worker devices across both scale-up fabrics inside a server and scale-out networks between servers; reducing bytes transferred or overlapping transfers with useful work can shorten that exposed time.

1.5.1.3 Orchestration coordination layer

Because a synchronous job waits for its slowest required participant, a single straggler or node fault can stall thousands of healthy peers even when both local kernels and network links have spare capacity. Coordination (C_3, the logic) captures the schedulers, synchronization barriers, and checkpointing protocols that manage the reliability gap (equation 1.2) and bind separate, failing

machines into one distributed computer. Its objective is goodput, the rate of useful training progress after wasted and repeated work are accounted for. It is one of several measures that share the name, and section 9.2.4 tells them apart. Overlap where dependencies permit and bounded recovery rollback improve that rate.

1.5.2 Projection from D·A·M to C^3

When a model or dataset that fit on one server is split across a cluster, the single-machine D·A·M components do not disappear; each one generates a distinct pattern of fleet-level compute, communication, and coordination work. The single-machine D·A·M taxonomy and the fleet-scale C^3 taxonomy therefore operate in tandem. D·A·M describes the workload: the *nouns* of the system (the training data, the executed algorithm, the hardware substrate). The C^3 taxonomy describes the execution cost: the *verbs* of the fleet (computing the math, communicating the results, coordinating the state). The C^3 taxonomy projects that workload onto a distributed fleet.

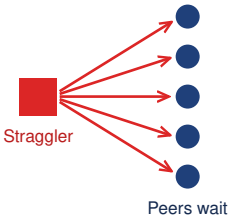

One straggler holds every synchronous peer.

Stretching an algorithm across nodes creates communication (C_2), such as the activation and gradient exchanges that tensor and pipeline parallelism add between partitions of one model. Stretching data across nodes also moves shards and gradients over the fabric (C_2), as when AllReduce synchronizes data-parallel replicas, and it requires coordination (C_3) through distributed samplers, gradient-aggregation barriers, and non-blocking checkpointing. Finally, managing hardware faults across a warehouse-scale machine requires orchestration and dynamic node recovery (C_3). Fleet-scale ML systems engineering requires solving the cross-products of D·A·M and C^3 rather than addressing each dimension in isolation.

1.5.3 The C^3 intersection landscape

The dimensions interact at their pairwise intersections, where an intervention along one axis alters costs across another. At the intersection of compute and communication ($C_1 \cap C_2$), overlapping transfers with arithmetic can hide network latency, but only while independent computation remains available. Between compute and coordination ($C_1 \cap C_3$), checkpointing preserves training progress against failure, yet it consumes execution resources; its frequency and placement trade present overhead against future recovery loss. Across communication and coordination ($C_2 \cap C_3$), collective synchronization depends jointly on transfer completion and participant readiness, so faster physical links alone cannot eliminate straggler waiting.

At the center ($C_1 \cap C_2 \cap C_3$), all three constraints converge simultaneously. Partitioning a model determines both which device computes each operation and what information must cross the network; recovering a failed partition requires restoring the corresponding shared state. These dependencies explain why parallelism, state sharding, and serving policies must be designed together. Each dimension contributes one term to the distributed step time, and each pairwise overlap names an interface where the dimensions trade costs by hiding latency, preserving progress against failure, and synchronizing collectives (figure 1.14).

The diagram organizes competing explanations for the same observed delay. A diagnosis may begin with one dimension or with a coupled mechanism, such as checkpoint traffic delaying a collective. The evidence determines which interaction to investigate; the diagram does not prescribe an order in which the dimensions must be optimized. Every chapter ahead uses the C^3 figure the same way, diagnosing which intersection binds before intervening. Surrounding these three operational dimensions are the non-negotiable operating invariants developed across Part IV, the Governance band of the fleet stack. They are the thermodynamic energy and cooling limits that constrain data centers, the adversarial security and differential privacy guarantees required of distributed pipelines, the robustness against silent corruption (a fault that changes a result without raising any error) and distribution shift (live inputs that move away from the training data), and the systemic governance, auditing, and safety invariants needed when operating at societal scale.

↪ Silent data corruption, and the checks that catch it across a fleet, are developed in section 7.5.

1.5.4 The Fleet Law

The C^3 intersection landscape defines *where* distributed execution spends time and where trade-offs bite. The Fleet Law turns that diagnostic map into a first estimate of distributed step time. Its terms are a simplified accounting model, so a prediction must state how work is divided and which costs overlap. On a single machine, execution time is governed by the iron law of performance,

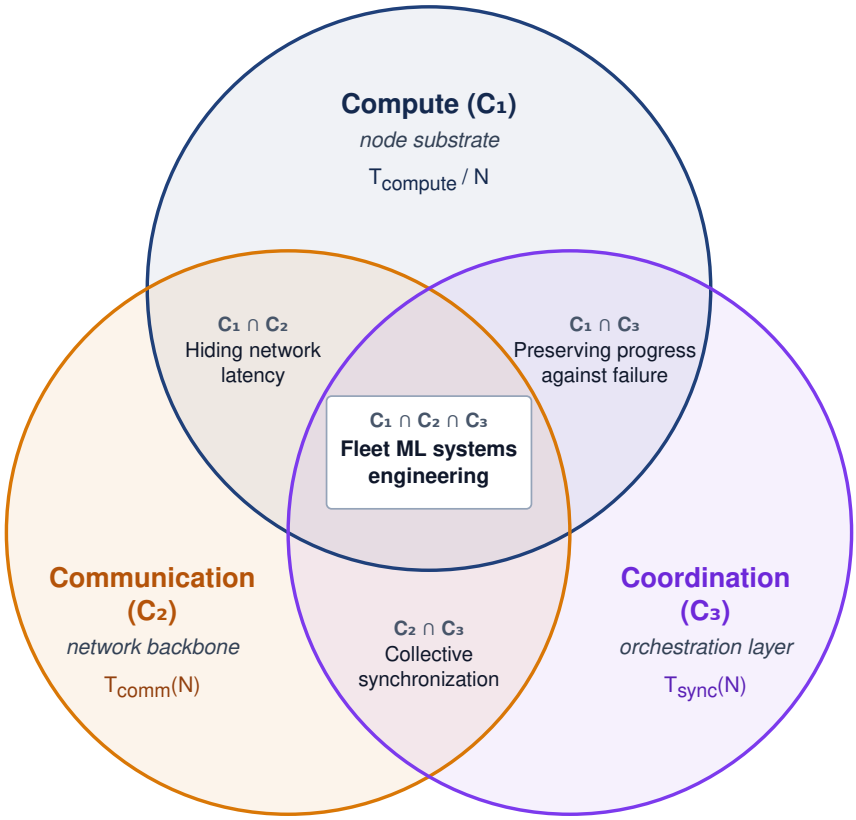

Figure 1.14: The C^3 Intersection Landscape: Each circle is one operational dimension (Compute, C_1, the node substrate; Communication, C_2, the network backbone; Coordination, C_3, the orchestration layer), labeled with the Fleet Law term it contributes to step time. Each pairwise lens names the interaction between two dimensions: hiding network latency behind computation, preserving training progress against failure, and synchronizing collectives. At the center, where all three meet, is fleet ML systems engineering. The techniques and chapters for each region are collected in table B.3.

$T = D_{\text{vol}}/\text{BW} + O/(R_{\text{peak}} \cdot \eta_{\text{hw}}) + L_{\text{lat}}$, decomposing into data movement, computation, and overhead. At fleet scale, that single-machine view no longer suffices. Equation 1.3 establishes the **Fleet Law**, which decomposes the idealized distributed training step into the terms of the C^3 engine:

$$T_{\text{step}}(N) = \frac{T_{\text{compute}}}{N} + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}} \tag{1.3}$$

Here T_{compute} is the compute term of the iron law above, the time one device spends on the forward and backward passes of one training step. Each term in the Fleet Law maps directly onto a region of the C^3 landscape in figure 1.14:

- $\frac{T_{\text{compute}}}{N}$ is the Compute (C_1) term, assuming ideally divisible local arithmetic for a fixed workload distributed across N devices.
- $T_{\text{comm}}(N)$ is the Communication (C_2) term, capturing the wire transit time moving gradients, activations, and parameter shards across the network fabric.
- $T_{\text{sync}}(N)$ is the Coordination (C_3) term, capturing time consumed by barrier semantics and straggler waiting, plus each step's average share of run-level costs such as cluster scheduling decisions and failure recovery.

- T_{overlap} is the Overlap ($C_1 \cap C_2$, $C_1 \cap C_3$, $C_2 \cap C_3$) term, representing communication or coordination hidden behind useful compute via techniques such as interleaving forward and backward work across pipeline stages, running transfers on separate execution streams, and writing checkpoints asynchronously.

The Compute term has two readings, and an analysis must state which one it uses. The arithmetic idealization used below assumes that local memory can supply operands fast enough to sustain the modeled compute rate, so the term holds arithmetic alone. A measured local-execution reference can instead include local memory stalls and dispatch costs in the Compute term, provided those costs divide with the work as assumed.

Doubling the device count shows why the subtraction matters. The local compute term falls to $T_{\text{compute}}/(2N)$, but communication and barrier synchronization persist, and the shorter compute window leaves less room to hide transfers, so T_{overlap} shrinks as well. The step on $2N$ devices therefore takes well over half as long as the step on N (figure 1.15).

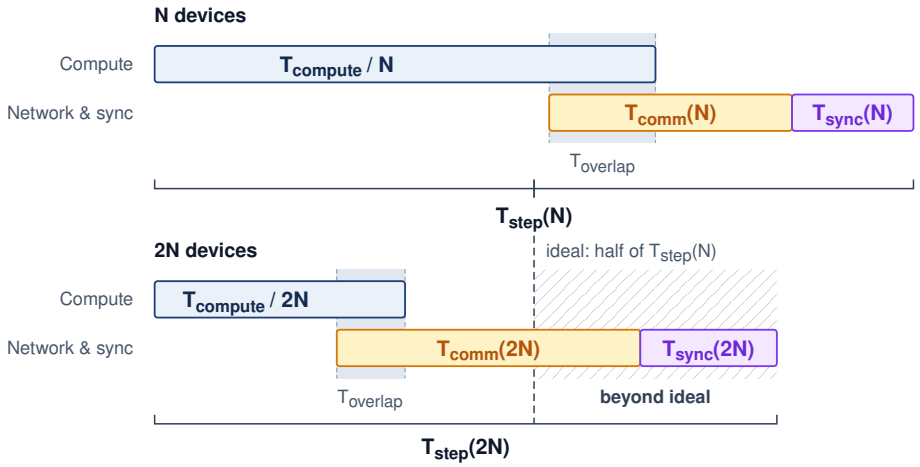

Figure 1.15: Anatomy of a Fleet Law Step: Each row is one training step for the same workload, with local computation on the upper lane and communication and synchronization on the lower lane. Read the step length from the bracket, not from any single bar: overlap (gray band) removes only the transfer time that runs beneath compute, so doubling the devices halves T_{compute}/N yet leaves $T_{\text{step}}(2N)$ well past the ideal half-length step (dashed line), by the hatched span. Bar lengths are illustrative, not measured.

Weak scaling (growing the workload with the device count), load imbalance, memory-capacity changes, and algorithmic changes require modifying these terms rather than treating T_{compute}/N as invariant. Measurements must also separate transfer from coordination time: a barrier interval that includes an unfinished transfer cannot be charged in full to both terms. Subtract only overlap supported by the execution timeline. Equation 1.4 rewrites this decomposition as the compute-time fraction. It measures useful arithmetic under the arithmetic idealization and local execution under the measured-reference convention:

$$f_{\text{compute}} = \frac{T_{\text{compute}}/N}{T_{\text{step}}(N)} \tag{1.4}$$

This decomposition is a diagnostic instrument. When f_{compute} drops below 0.5, more than half the step lies outside the modeled Compute term; under the arithmetic idealization, that means outside useful local arithmetic. The C^3 taxonomy identifies where the time goes. If $T_{\text{comm}}(N)$ dominates, upgrade interconnects or overlap communication with computation. If $T_{\text{sync}}(N)$ dominates, consider asynchronous methods or reduce barrier frequency. If T_{compute}/N dominates, the modeled local execution sets the step time. The **displacement of overhead** is a diagnostic heuristic: reducing one overhead can relocate cost among the three dimensions. Asynchronous training removes some

coordination barriers but introduces gradient staleness, in which workers apply updates computed from an older copy of the model. Pipeline parallelism can reduce some communication exposure but adds bubble time, the intervals in which pipeline stages sit idle waiting for work.

The compute-time fraction describes how a single N-device configuration divides its own step time. A production team also needs a scalar that compares that distributed step against a single-device baseline to show whether adding devices actually speeds up the job. The scaling efficiency (η_{scaling}) of the fleet is the ratio of ideal parallel step time to actual step time:

$$\eta_{\text{scaling}} = \frac{T_1}{N \times T_N}$$

where T_1 is the single-device time for the same work and $T_N = T_{\text{step}}(N)$ is the step time on N devices. When local execution divides cleanly by N, T_1/N matches T_{compute}/N and the two ratios agree. In practice, partitioning a workload across more devices can also shrink per-device batch slices enough to lower local kernel utilization or add recomputation work, which reduces η_{scaling} even before communication and synchronization time are counted. Section 6.4.1 develops the measure in full.

↪ The component decomposition that maps each term of $T_{\text{step}}(N)$ to measurable activities is developed in section B.3.1.

> **Systems Perspective 1.1: Amdahl's distributed pitfall**
>
> The Fleet Law applies Amdahl's reasoning (Amdahl 1967) to a strong-scaling experiment: a fixed amount of useful work is divided across more devices while communication and coordination terms are measured for that configuration. A common shortcut reads a measured 20 percent network share as a 5× speedup ceiling, but Amdahl's bound applies only if that 20 percent of the original workload remains irreducibly serial as compute is accelerated. A network share measured at one device count is not automatically invariant; changing the partition, topology, collective algorithm, or overlap can change both its absolute time and its fraction of the step. The Fleet Law therefore diagnoses the scaling limit of a stated workload and configuration rather than asserting a universal network fraction.

Fleet efficiency has several measures, and each answers a different question:

- Percent of peak compares achieved arithmetic throughput with peak throughput over a stated measurement interval, and the result depends on what it counts as arithmetic.
- MFU is percent of peak that counts only the model's own forward and backward work, and GPT-3's roughly 20 percent is a measure of this kind. GPT-3's estimate averages over the training run, so it includes idle time within that interval; a measurement restricted to active kernels uses a different denominator.
- Hardware FLOPs utilization is percent of peak that counts every executed operation, including activations discarded to save memory and recomputed during the backward pass. Megatron-LM's 52 percent is a measure of this kind, so it runs higher than the MFU of comparable runs, which the PaLM report puts at 30.2 percent for Megatron-Turing NLG and 46.2 percent for PaLM (Chowdhery et al. 2022).
- The compute-time fraction asks how much of one step belongs to the modeled local-execution term.
- Scaling efficiency asks how close N devices come to an N-fold speedup over one.
- Goodput widens the window from one step to the whole run by discounting work lost to failures and repeated after restarts; Meta's report of more than 90 percent effective training time for Llama 3 is a measure of this kind.
- Google's ML productivity goodput shares the name but answers a fleet-wide question, the fraction of capacity that becomes useful work, as the product of scheduling, runtime, and program goodput (Wongpanich et al. 2025). Section 9.2 defines the three factors and the boundary at which each one counts a loss.

A fleet can score well on one measure and poorly on another, so a claim about fleet efficiency must name the measure and measurement interval it uses.

The Fleet Law, the compute-time fraction, and scaling efficiency together describe the time side of fleet efficiency. They separate modeled local execution from exposed communication and synchronization costs and compare the resulting step time with the single-device baseline. Fleet efficiency also has an energy side, because every exposed transfer and every idle accelerator still draws power.

1.5.5 The energy accounting boundary

While the Fleet Law accounts for time, comparing fleet designs on energy requires a fixed **energy accounting boundary**, the set of energy consumers that every design in the comparison is charged for. The boundary is what stays fixed, not any particular efficiency value. Equation 1.5 defines **fleet energy productivity** (ρ_{energy}), measured in FLOP/J, within a simplified operational boundary:

$$\rho_{\text{energy}} = \frac{O_{\text{useful}}}{E_{\text{device}} + E_{\text{cooling}} + E_{\text{network}}} \tag{1.5}$$

where O_{useful} is useful work in FLOPs. The three energy terms cover compute devices, attributable cooling, and the network fabric, without counting the same consumption twice. Host, storage, and other facility loads must be added explicitly for a complete operational estimate; embodied impacts, the environmental costs of manufacturing the equipment, lie outside this boundary. Comparisons must use the same boundary and useful-work definition. The resulting Pareto frontier balances minimizing T_{step} with maximizing ρ_{energy}.

The 175B reference job connects the time and energy accounts, because every second of exposed communication or pipeline idle time both delays the step and keeps accelerators and links drawing power. Notebook 1.1 computes both costs for that job on 1,024 and on 8,192 GPUs.

↪ Power usage effectiveness, the ratio of total facility energy (including cooling and power conversion) to IT energy, and what it leaves out are developed in section 15.4.1.3.

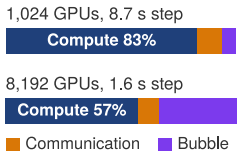

Spreading the same batch over 8× the GPUs cuts compute to 57 percent of the modeled step, from 83 percent.

Napkin Math 1.1: The time and energy cost of scaling out

Problem: The 175B reference job runs on 1,024 H100 GPUs as 16 replicas of 64 GPUs, with a global batch of 1,536 sequences per step. On 8,192 GPUs, the same batch spreads over 128 replicas of 12 sequences each instead of 96. How do the modeled compute-time fraction and the energy of one step change between the two clusters?

Physics: The step model used here divides each step into the replica's local computation, the tensor-parallel exchanges inside each node, the data-parallel gradient exchange across the fabric, and the pipeline bubble, the time stages sit idle while the pipeline fills and drains. The gradient exchange runs during the backward pass, and the model hides 85 percent of it behind that computation, leaving the rest exposed. These are modeled values, not measurements.

Analysis: On 1,024 GPUs, local computation takes 7.25 s of the 8.7 s step. Tensor-parallel exchanges add 0.90 s and the bubble 0.53 s, while only 0.03 s of the 0.21 s gradient exchange is exposed. The compute-time fraction is 83 percent. On 8,192 GPUs, local computation falls to 0.91 s and tensor-parallel traffic to 0.11 s, because each replica carries fewer sequences. The gradient exchange stays near 0.22 s, and the bubble does not shrink at all, because it is set by the time one sequence spends in each pipeline stage. With fewer sequences to keep the pipeline busy, that fixed idle time fills 33 percent of the step. The step shrinks only to 1.6 s, a 5.5× speedup for 8× the GPUs, and the compute-time fraction falls to 57 percent. Because the compute term is the replica's whole local workload, this fraction is also the step's scaling efficiency, 0.83 and 0.57.

Energy: Both clusters perform the same optimizer update, so the useful FLOPs in equation 1.5 are identical. Charging each transmitted bit once at 35 pJ/bit, the gradient exchange costs about 2.9 kJ of fabric energy per step on the smaller cluster and about 24.9 kJ on the larger, because 8× as many GPUs each send nearly the same gradient slice. The accelerators matter more. The step occupies 8,921 GPU-seconds on the smaller cluster and 12,983 on the larger, and an accelerator idle in the bubble or waiting on an exchange still draws power. If each accelerator draws similar power whether it computes, exchanges, or waits in the bubble, fleet energy per step grows with those GPU-seconds, so energy productivity falls as the fleet grows, even though each step finishes sooner.

Systems insight: The model prices links at line rate, hides a fixed share of the gradient exchange, treats tensor-parallel exchanges as unoverlapped, and ignores stragglers, so a measured step at either scale can differ. The trend is the lesson. Adding accelerators to a fixed batch shrinks the compute term in proportion, while the bubble and the gradient exchange do not shrink, so exposed communication and pipeline idle time claim a growing share of each step. Section 1.2.4 found the same pressure on the network side, where the gradient exchange climbs from 6 percent to 53 percent of fabric balance.

↪ Why the pipeline bubble grows as each replica receives fewer microbatches is developed in section 6.5.2.

↪ The energy that exposed collectives and idle, synchronization-stalled accelerators add to a distributed training run is measured in section 15.7.2.1.

1.6 The Fleet Stack and the Lighthouse Archetypes

When a distributed step stalls on exposed communication or synchronization, the Fleet Law and the C^3 dimensions tell us which cost dominates, but not which layer of the system caused it or which workloads will trigger it next. Three architectural maps work together to close that gap: the fleet stack locates system responsibility, C^3 and the Fleet Law diagnose execution cost, and the lighthouse archetypes test those costs across representative production workloads. The fleet stack, traced upward from the hardware of one node in section 1.2, arranges the fleet into eight rows grouped into four ascending bands (figure 1.4). From the bottom, those rows are Facility and Accelerators, Fabric and Storage, Distributed Execution, Reliability and Control, Runtime and Serving, Operations, Assurance, and Governance and Sustainability. Capabilities fixed in the lower rows, from facility cooling and RDMA fabrics up through distributed execution, constrain the serving runtimes and governance mechanisms above them. Those four bands correspond to the book's four parts: the physical fleet, distributed ML, the fleet in production, and the responsible fleet. The chapter map at section 1.7 follows that dependency order.

Six **engineering practices** guide decisions at every band of the stack: instrument first, reserve headroom for recovery and tail latency, expect bottlenecks to migrate, treat failure as steady state, compound small efficiency gains across the fleet, and co-design hardware with algorithms. For example, faster accelerators can expose a network bottleneck; instrumentation identifies that migration, and hardware co-design aligns the fabric with the workload. The practices describe how a fleet engineer works. The principles of distributed ML systems, among them that communication can dominate a step and that failure is routine, describe what the physics of the fleet makes true, whatever the practice. Each part of this book establishes its share of those principles. Section 17.2 gathers them into six, one or two for each part, and pairs each with the practice it demands.

1.6.1 Three systems archetypes

A single benchmark workload cannot test every band of the fleet stack, because a dense language model, a sparse recommendation service, and an on-device federated learner fail in completely different places. Three lighthouse archetypes at scale recur across the book where they illuminate a chapter's central constraint; no chapter must force all three into its argument. Together they sample distinct C^3 regimes by stressing communication, coordination, or compute in different ways.

1.6.1.1 Archetype A (GPT-4/Llama-3)

The training state of a dense language model with hundreds of billions of parameters cannot fit in the accelerator memory of a single server. Splitting the model across thousands of accelerators using 3D parallelism (data, tensor, and pipeline) risks exposing enough gradient and activation traffic to make the network the binding bottleneck. GPT-4 and Llama 3, the autoregressive transformer architectures behind this archetype, generate text one token at a time. At single-accelerator scale, autoregressive language models are memory-bandwidth probes. At the scale of hundreds of billions of parameters, and in some proprietary systems potentially more, training demands EFLOP/s of sustained compute distributed across thousands to tens of thousands of accelerators. Serving requires memory bandwidth sufficient to load billions of weights for every generated token. Because gradient and activation traffic crosses the network every step, Archetype A probes the communication dimension of C^3, although its actual dominant term depends on the model, partition, topology, and degree of overlap.

Where a chapter needs a concrete instance, Archetype A is a 175B-parameter dense transformer of the GPT-3 class, chosen because its configuration is published while GPT-4's is not, trained with mixed precision and the Adaptive Moment Estimation (Adam) optimizer. Its training state occupies 16 bytes per parameter (BF16 weights and gradients, FP32 master weights, and two FP32 Adam moments), or 2.8 TB. A resumable checkpoint omits the gradients and needs 14 bytes per parameter, or 2.45 TB. The reference cluster is 1,024 H100 GPUs, arranged as 128 nodes of 8 accelerators each. Each data-parallel worker is one replica of the model spread over 64 GPUs, with tensor parallelism across the 8 GPUs of a node and a pipeline 8 nodes deep, so the cluster holds 16 replicas. Without sharding, the training state then comes to 43.8 GB per GPU, which fits in one H100's memory. One training step is one optimizer update over a global batch of 1,536 sequences (3,145,728 tokens), with one gradient reduction across the replicas, and the modeled step takes about 8.7 s. A chapter that studies a larger cluster or a different model says explicitly how it varies this reference.

1.6.1.2 Archetype B (DLRM at Scale)

Where a dense language model spends most of its memory on weights that every token uses, a production recommender spends its memory on sparse embedding tables that reach 10 TB or more across hundreds of nodes. It serves roughly 10 million queries per second (QPS) within a <100 ms tail-latency budget. DLRM, the deep learning recommendation model that Meta published as a reference architecture for personalized content ranking, derives its capacity from dense matrix multiplications, as LLMs do, and also from large embedding tables that map billions of user and item features into dense vectors. Those tables far exceed any single accelerator's memory, so they are sharded across devices while each worker processes its own slice of the request batch for the dense layers.

Because a worker's requests reference embedding rows stored on many other devices, serving and training must coordinate AllToAll exchanges, collective operations in which every worker sends a distinct subset of looked-up vectors to every other worker across $\mathcal{O}(N^2)$ potential communicating pairs. The resulting byte volume and contention depend on the request batch, partitioning, and network topology. In the C^3 taxonomy, the AllToAll bytes are measured as Communication, but Archetype B stresses coordination, because sparse feature routing, shard placement, request scheduling, and embedding-update consistency determine whether that traffic can be served within tail-latency limits.

1.6.1.3 Archetype C (Federated MobileNet)

Both of those workloads run inside managed data centers. Federated learning instead trains a model across millions of varied smartphone-class devices whose power budgets span roughly 2 W to 10 W (a few watts sustained, with brief bursts at the top of that range), without ever moving raw user data off the device. MobileNet, a family of efficient convolutional neural networks designed for on-device execution under tight memory and power budgets, serves as the reference workload for this federated learning setting. Microcontroller-class sensors that run on hundreds of kilobytes of memory and milliwatts of power lie at an extreme below this archetype.

In the C^3 taxonomy, Archetype C probes the per-device compute envelope: local execution must fit within the device's power, memory, and thermal limits. That local constraint does not by itself determine the bottleneck of a federated round. Uploading updates can expose communication limits, while selecting available participants and waiting for their results can expose coordination limits. The fleet challenge is coordinating model updates across millions of such compute-constrained, unreliable devices using federated averaging, which periodically averages device updates at a coordinator. Data distributions that are not independent and identically distributed (i.i.d.) and intermittent participation affect the resulting learning process; convergence claims must state the assumptions under which they hold.

The three archetypes are constraint probes, not a catalog of model families. The LLM case asks whether dense synchronization can keep thousands of accelerators useful, the DLRM case asks whether sparse feature routing can meet tail-latency budgets, and the federated MobileNet case asks whether local device constraints can be coordinated into a statistically coherent update. Each of the three canonical workloads pairs a C^3 diagnostic regime with the fleet challenge it raises (table 1.3), and section 1.7 traces how the chapters climb the fleet stack to meet them.

Table 1.3: Lighthouse Archetypes at Scale: The three canonical workloads in representative diagnostic regimes, with the fleet-scale challenge each raises. These are constraint probes, not invariant bottleneck assignments; workload, partitioning, topology, and overlap determine the binding cost. The MobileNet row describes the local device envelope, not the bottleneck of a complete federated round.

Archetype	C^3 diagnostic regime	Fleet challenge
Archetype A (GPT-4/Llama-3)	Communication (fleet-wide)	Partition hundreds of billions of parameters, and potentially larger proprietary models, across thousands to tens of thousands of GPUs using 3D parallelism without the network becoming the bottleneck.
Archetype B (DLRM at Scale)	Coordination (routing and placement)	Shard 10 TB+ embedding tables across hundreds of nodes; process millions of QPS with <100 ms tail latency while managing sparse feature routing, shard placement, and exchanges among $\mathcal{O}(N^2)$ potential pairs.
Archetype C (Federated MobileNet)	Compute (per-device envelope)	Coordinate learning across millions of compute-constrained, unreliable edge devices using federated updates; raw data cannot leave the device.

1.7 The Structure of This Textbook

Meeting those three fleet challenges requires climbing the fleet stack in dependency order, because choices fixed in the physical substrate constrain every algorithm, runtime, and governance control built above it. The chapter sequence is therefore a dependency map rather than a topic inventory. Silicon, cooling, and network limits constrain training partitioning; partitioned algorithms create fault-tolerance and orchestration obligations; and those mechanisms support production operation. Requirements also flow in the other direction: security, governance, and service targets constrain which lower-layer designs are acceptable. Part I builds the physical cluster substrate, Part II coordinates distributed execution and fault recovery, Part III runs the fleet in production, from efficiency and serving through edge deployment and operations, and Part IV enforces governance, security, and sustainability invariants (figure 1.16).

The infrastructure arc begins with a limit no algorithm can ignore, because no single server has enough memory, power, cooling, or I/O to train the largest models. Chapter 2 therefore starts with high-density silicon and facility physics, chapter 3 turns isolated machines into a cluster-scale gradient bus, and chapter 4 completes the substrate by keeping accelerators fed without letting storage become the hidden bottleneck.

The distribution arc begins once the physical fleet exists and coordination costs become unavoidable. Splitting math across machines creates two new problems: replicas must exchange state fast enough to behave like one optimizer, and failures become routine rather than exceptional. Chapter 5 develops the communication primitives that bind independent nodes into a coherent computer and prices each exchange. Chapter 6 uses them to build partitioning strategies for models with hundreds of billions to trillions of parameters. Chapter 7 treats recovery speed as more important than individual-node uptime. Chapter 8 manages the multi-tenant cluster where those jobs must be placed, isolated, and rescheduled.

The production arc puts the built fleet to work on real jobs and live traffic. Once a fleet runs in production, the binding question is no longer only time-to-train; it is whether utilization, latency, and lifetime serving cost stay inside the operating envelope. Chapter 9 measures how much of the fleet's capacity becomes useful work, in a training step and a serving replica alike, and charges each loss to scheduling, runtime, or the program. Chapter 10 turns each replica's measured capacity into admission, batching, and routing policy for millions of users. Chapter 11 moves learning onto phone-class devices whose memory, battery, and thermal budgets decide what adaptation they can afford. Chapter 12 develops the shared controls (registries, release gates, monitoring, and cost accounting) that govern hundreds of models as one fleet.

The governance arc treats responsibility as an engineering layer rather than an afterword. At global scale, technical bugs become societal hazards: adversaries can poison data or extract weights, open-world inputs can break brittle models, energy use becomes a fleet constraint, and governance

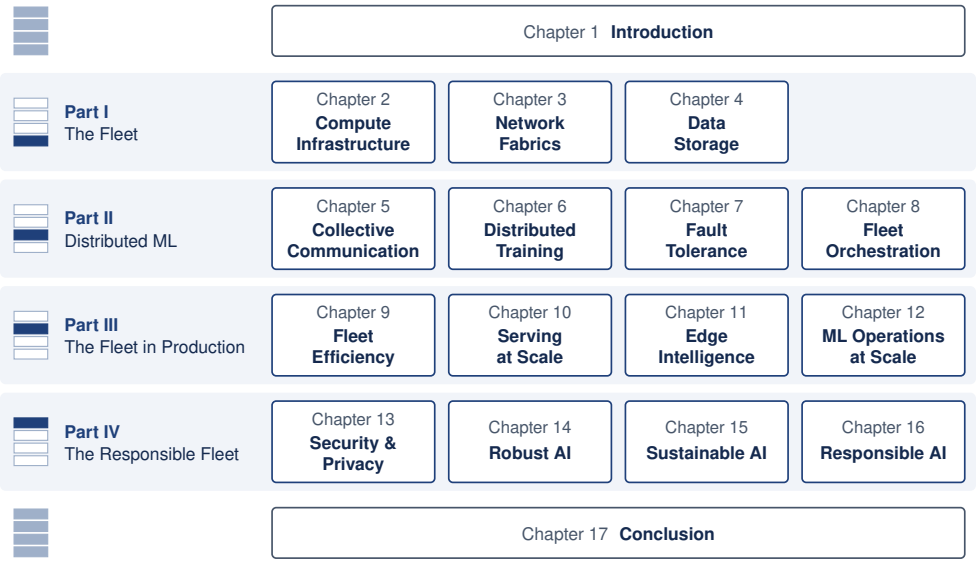

Figure 1.16: The Chapter Roadmap: The chapters in reading order, top to bottom. The introduction frames the book; the four parts then climb the fleet stack of figure 1.4 from physical infrastructure through distribution and production to governance, as the small stack beside each part shows by lighting that part's band; the conclusion gathers the principles of all four parts.

failures can turn technical capability into social harm. Chapter 13, chapter 14, chapter 15, and chapter 16 develop those obligations as system properties that must be designed, measured, and operated.

1.8 Fallacies and Pitfalls

The following fallacies and pitfalls capture architectural mistakes that waste development resources, miss performance targets, or deploy systems critically mismatched to their operating constraints. Each represents a pattern observed repeatedly in the transition from single-machine ML to the machine learning fleet.

Fallacy: *Accelerators twice as fast halve the training step.*

An accelerator upgrade that changes only arithmetic throughput shrinks the T_{compute}/N term of the Fleet Law while leaving the communication terms unchanged. On the 1,024-GPU reference job of notebook 1.1, local computation takes 7.25 s of the 8.7 s modeled step, and the pipeline bubble, idle time measured in stage computation, adds 0.53 s. Accelerators 2× faster halve both, but the 0.93 s of exposed tensor-parallel and gradient exchange does not speed up, so the step falls only to about 4.8 s, a 1.8× speedup rather than 2×. The larger the share of the step that communication occupies, the further the gain falls short. The direction of the effect is general: relieving the Compute axis tends to expose Communication or Coordination as the binding cost, the displacement of overhead described with the Fleet Law. A real upgrade may also change memory capacity, transfer rates, or available overlap. An accelerator purchase must therefore be judged against the fleet's measured step decomposition, not against the datasheet ratio.

Pitfall: *Assuming that reliable components make a reliable job.*

A per-GPU MTTF of 50,000 hours describes a part that rarely fails. A synchronous job, however, needs all of its parts at once. As section 1.1 showed, a job spanning 25,000 GPUs has an MTBF of about 2 hours, or roughly 12 interruptions per day. Component reliability characterizes one part; the reliability gap (section 1.4.1) characterizes the job, whose availability decays exponentially with the number of parts it requires simultaneously. Recovery practices that suit a single machine, such

as manual restart and infrequent checkpoints, fail at fleet scale even though no individual part has become less reliable.

Fallacy: *Edge deployment efficiency requirements are simply scaled-down versions of cloud requirements.*

The fallacy treats edge systems as resource-constrained cloud systems. Edge devices face qualitatively different constraints. For example, an autonomous vehicle traveling at 120 km/h covers 3.33 m during every 100 ms of processing delay, adding that distance to its reaction distance. While cloud deployments scale to kilowatts, a phone-class device of the kind Archetype C models sustains only a few watts. A cloud-optimized model with 95 percent accuracy and 50 ms latency might be unusable on an edge device where thermal throttling increases latency to 200 ms and drains the battery in under an hour.

Pitfall: *Treating a compute budget as a wall-clock performance prediction.*

This pitfall, sometimes called the linear scaling fallacy, confuses useful computational work with the time required to execute it. The empirical loss relation $\mathcal{L}(X) = A_{\mathcal{L}} X^{-\alpha_{\text{scale}}} + C_{\mathcal{L}}$ estimates model quality within a validated regime; the Fleet Law separately estimates the time needed to perform that work. A team distributing a fixed training workload might predict a 3× improvement but achieve only 1.8× if coordination and communication overhead consumes 40 percent of the scaled step time. This shortfall does not by itself invalidate a loss scaling law. It shows why a resource plan needs both a model-quality estimate and a measured execution-cost estimate.

1.9 Summary

Fleet planning now separates two questions that single-machine practice could merge. Scaling laws answer how much model size, data, and useful computation a quality target requires; the Fleet Law answers how long a fleet takes to execute that work. Single-machine diagnostic skills remain necessary at that scale, but the engineering boundary expands, because a fleet can lose useful progress to network transfers, inconsistent shared state, or failed participants even when its local kernels are efficient. C^3 distinguishes compute, communication, and coordination costs, while the fleet stack locates the decisions that shape them. With these tools, a resource plan can be traced through to execution behavior, subject to the reliability, energy, and governance constraints of the mission.

Key Takeaways: Scale makes a new machine

- **More devices change the engineering boundary**: Distributing work introduces shared network and failure dependencies. Synchronous training couples every worker to one shared step, so a single slow or failed participant holds back the whole job. Local performance remains necessary, but fleet decisions must also account for synchronization, recovery, power delivery, and governance across that wider boundary.
- **Communication can consume the speedup**: Dividing fixed work across accelerators shortens ideal local compute time while exposing communication and coordination costs. A 175B-parameter model can move hundreds of gigabytes of gradients per step. Comparing communication intensity, the bytes moved per local operation, with the fabric's balance point shows whether local computation or the network sets the step time.
- **Failure becomes steady state**: An unreplicated synchronous job needs every required component at once, so with independent component failures its availability falls exponentially with job size, to only 36.8 percent for 1,000 required nodes at an illustrative 99.9 percent each. In the modeled 25,000-GPU job, those failures interrupt tightly coupled training every few hours despite long per-GPU lifetimes, so checkpointing, recovery, and observability must be designed into normal operation.
- **Scaling laws set the demand; the Fleet Law sets the price**: Within a validated regime, loss falls as a power law in compute, data, and parameters, so each equal quality gain needs disproportionately more resources, and compute-optimal plans grow parameters and tokens together. Those laws say how much work is worth buying; the Fleet Law says

how long the fleet takes to deliver it. Task-level evaluation, not loss alone, establishes whether the predicted gain justifies the resources, and training-compute optimality alone does not establish the lowest lifetime serving cost.

- **C^3 extends local intuition**: The Data, Algorithm, Machine lens describes the workload and substrate; Compute, Communication, and Coordination expose the costs of fleet execution. The Fleet Law connects useful local work to transfer, synchronization, and overlap costs as device count changes, and the fleet stack locates the layer where each decision belongs.
- **Impact needs a control plane**: Foundation-scale systems amplify security, privacy, fairness, and policy failures across large user populations. Governance is not an appendix to infrastructure; it is the mechanism that decides what the fleet is permitted to optimize and deploy.

The premise of this book is that scale can move the binding constraint outside the node. A local speedup is therefore evidence about one component, not yet evidence that the fleet makes more useful progress. Engineering judgment requires tracing that change through communication, coordination, and the mission's operating constraints before accepting it as an improvement.

What's Next: From requirements to silicon

The scaling laws set how much work a frontier model demands, the walls of the scale moment mark where that demand meets physical limits, and the Fleet Law prices each distributed step. Those requirements say what the fleet must sustain, not what hardware can sustain it. Chapter 2 turns to that hardware. It begins where the scaling laws leave a growing model, too large for one accelerator's memory, and climbs from the accelerator through the node and rack to the pod. At each scale it asks which physical resource runs out first: memory bandwidth, power and cooling, communication bandwidth, or failure exposure. Those limits form the Facility and Accelerators row at the base of the fleet stack, the substrate on which every later chapter builds.

I

THE FLEET

Part I

1

Fleet Principles

The scaling laws of section 1.3 reward more parameters, more training tokens, and more compute together, so each further gain in model quality asks the fleet for more memory, more arithmetic, more bytes moved between machines, and more data stored and read. With that demand set, Part I builds the physics of supply: the silicon, the wires, the cooling systems, and the storage hierarchies of the warehouse-scale machine learning fleet, where the network is the bus, power density limits deployment, and component failures are routine at scale. The broader fleet also includes geographically distributed serving systems and heterogeneous edge devices, whose constraints appear in Part III. Across this physical foundation, the focus shifts from the single accelerator to the data-center-scale machine, and the engineering question is no longer how to compute on one device but how to move energy and information at a scale that challenges the limits of the physical infrastructure.

At this scale the individual accelerator becomes one component of a larger, tightly coupled system. The principles of the fleet are therefore not best practices for cluster management. They are constraints, most of them physical, that dictate what kind of models can be trained and how they can be served. These constraints range from the efficiency of specialized silicon and the heat a rack can shed to the bisection bandwidth of the network fabric. They define the boundaries of the Infrastructure band of the fleet stack, the two rows (Facility and Accelerators, Fabric and Storage) on which every other part of the book rests. The C^3 taxonomy tracks costs across Compute, Communication, and Coordination. An optimization can reduce total execution time while shifting which cost dominates; reducing local compute, for example, can expose communication that was previously hidden. The engineering obligation is to measure the complete path rather than assume that a local saving yields the same fleet-wide gain.

The principles below preview the constraints that Part I makes concrete through compute infrastructure, network fabrics, and storage, in the order those chapters develop them. Their implications name some mechanisms developed later in the book; those names are signposts, not prerequisites for reading this part. The first constraint applies before any machine is assembled, because the choice of accelerator decides how much of its silicon and power go to useful arithmetic.

Principle 1: The generality tax

Invariant: For workloads with stable computation patterns, specialization can improve performance per watt by reducing general-purpose overhead, while also reducing programmability and workload flexibility.

$$\rho_{\text{energy,specific}} \gg \rho_{\text{energy,general}} \quad \text{for a fixed power budget}$$

Here, ρ_{energy} denotes useful computation per unit energy; the comparison illustrates the potential advantage for a workload suited to the specialized design.

Implication: The trajectory from CPU to GPU to Tensor Processing Unit (TPU) to fixed-function ASIC is a workload-dependent architecture trade-off, not a universal physical law. Each step can trade programmability for efficiency, and the fleet architect must choose the right point on this curve for each workload.

A specialized accelerator still holds only as much model state as its memory allows, and that capacity gives the first reason to divide a model. When parameters, gradients, and optimizer states outgrow a single package, the computation can proceed only by splitting that state across multiple devices.

Principle 2: The memory capacity gap

Invariant: Frontier model memory requirements have grown faster and more irregularly than single-accelerator HBM capacity.

Implication: When a model and its training state exceed one device's memory, the system must divide that state across devices. Part II develops the parallelism strategies for doing so: tensor parallelism divides individual operations, and pipeline parallelism assigns groups of layers to different devices. Complementary data parallelism distributes training examples across model replicas. Because per-device bandwidth falls roughly tenfold at each physical boundary, from device memory to the node interconnect to the cluster fabric, the division that exchanges the most data, tensor parallelism, stays inside the fastest domain. Part I establishes the capacity and communication constraints those strategies must respect.

Dividing a model across more accelerators packs more of them into each rack, and the rack's power budget is the first limit on that density. Keeping latency-sensitive peers close enough for high-bandwidth copper links concentrates electrical draw into a narrow physical footprint.

Principle 3: The power density wall

Invariant: Cluster scaling is constrained by thermal dissipation limits (watts per rack) and cooling capacity, not just silicon area or floor space.

Implication: Watts per rack, together with power-delivery losses and the cooling class the facility installs, bound how many accelerators can share a rack, and therefore how large a fast interconnect domain can be. Accelerator power has risen faster than rack power budgets, so installable compute is set by the rack envelope rather than by peak FLOP/s.

Every watt the rack draws becomes heat that the facility must remove. Sustaining that electrical load therefore turns a power-delivery specification into a continuous heat-rejection requirement.

Principle 4: The thermodynamic limit

Invariant: Irreversible computation and data movement dissipate electrical energy as heat; at data-center scale, the practical limit is the rate at which that heat can be removed.

$$\dot{Q}_{\text{IT}} \leq \dot{Q}_{\text{remove,max}}(\text{facility}, \text{cooling configuration})$$

where $\dot{Q}_{\text{IT}}$ is the heat produced by IT equipment and $\dot{Q}_{\text{remove,max}}$ is the maximum heat-removal rate of the installed facility and cooling configuration.

Implication: A rack can operate only while its heat output remains within the installed cooling system's capacity. The cooling class is chosen to match that heat load. Engineered air cooling remains viable for configurations designed for its airflow, while denser configurations need cold-plate or immersion liquid loops, which carry the same heat with far less flow at a cost in serviceability and leak and loss-of-flow risk.

Once a model is divided across devices, the network becomes part of the execution path. The time to move a single message has a structure of its own, which determines whether a transfer is limited by fixed startup cost or by bandwidth.

Principle 5: The bandwidth-latency trade-off (α-β model)

Invariant: In the latency-bandwidth model, communication time combines fixed startup latency (α) with the transfer time for a message of size n at effective bandwidth β.

$$T(n) = \alpha + \frac{n}{\beta}$$

Implication: Small messages (for example, routing metadata or control signals) are latency-bound; large messages (for example, gradients) are bandwidth-bound. Optimization strategies must match the regime: fuse small messages to amortize α, compress large messages to reduce n.

The α-β model prices one message on one path. When every worker communicates at once, those messages compete for the links that join one half of the cluster to the other, and the capacity of that cut bounds their aggregate rate.

Principle 6: The bisection bandwidth theorem

Invariant: Communication that crosses a network cut is limited by the aggregate bandwidth of that cut. For traffic that must cross between two halves of a cluster, the minimum bisection bandwidth bounds the achievable transfer rate.

Implication: Synchronized training requires bisection bandwidth commensurate with the traffic crossing the fabric. A collective operation coordinates communication among a group of workers; AllReduce, for example, combines their contributions and returns the result to every worker. Part I explains how network topology limits this traffic; Part II develops the collective algorithms that generate it.

A fabric that carries every exchange on time still helps only if the training pipeline can keep the accelerators fed. Every training step begins by pulling fresh batches from storage, so an underprovisioned input path stalls the cluster before arithmetic or gradient exchange even starts.

Principle 7: The I/O wall

Invariant: When storage throughput cannot deliver training data as fast as accelerators consume it, accelerators idle regardless of their computational power.

$$\text{Utilization} = \min\left(1,\ \frac{\text{BW}_{\text{storage}}}{N_{\text{GPU}} \times R_{\text{consumption}}}\right)$$

where $\text{BW}_{\text{storage}}$ is aggregate storage bandwidth, N_{GPU} is the number of accelerators, and $R_{\text{consumption}}$ is each accelerator's data-consumption rate.

Implication: A storage system that was adequate for 8 accelerators can become the bottleneck at 64 if its provisioned throughput cannot meet the higher aggregate demand. The I/O wall scales with the number of accelerators: at a fixed per-device consumption rate, every accelerator added raises the throughput floor that storage must sustain. Without matching capacity, the data pipeline can become the limiting factor.

Meeting that demand requires a hierarchy rather than a single undifferentiated storage pool. No single storage medium offers both the petabyte capacity of a training corpus and the terabyte-per-second bandwidth of an accelerator pod.

Principle 8: The storage hierarchy principle

Invariant: Storage hierarchies are designed so that tiers farther from the accelerator typically trade higher access latency and lower per-stream bandwidth for greater capacity and lower cost per byte; aggregate throughput still depends on provisioning and parallelism.

Implication: The systems engineer's task is to ensure that the right data is on the right tier at the right time. Data format choices, caching strategies, prefetch buffer sizing, and tiering policies all exist to manage movement upward through the hierarchy so that accelerators do not starve.

These constraints make the fleet a first-class engineering object. Following the introduction's account of why single machines no longer suffice, Part I builds the physical machine from the ground up, starting with the silicon, power, and cooling systems of the AI data center and the network fabrics that connect the fleet. It then turns to the storage hierarchy that feeds the training pipeline, preserves its checkpoints, versions the artifacts a run depends on, and serves retrieval at inference time. Each chapter sets one or more terms of the Fleet Law (equation 1.3). Compute infrastructure sets the per-device rate behind T_{compute}/N and the failure exposure that later surfaces in $T_{\text{sync}}(N)$. Network fabrics set $T_{\text{comm}}(N)$ and the barrier tail. Storage adds the input stalls and checkpoint pauses that the step pays whenever prefetching and asynchronous writes cannot hide them. Together, these chapters fix the physical rates within which the communication, partitioning, and recovery mechanisms of Part II must work, for each of the lighthouse archetypes (section 1.6.1) traced across this book.

2

Compute Infrastructure

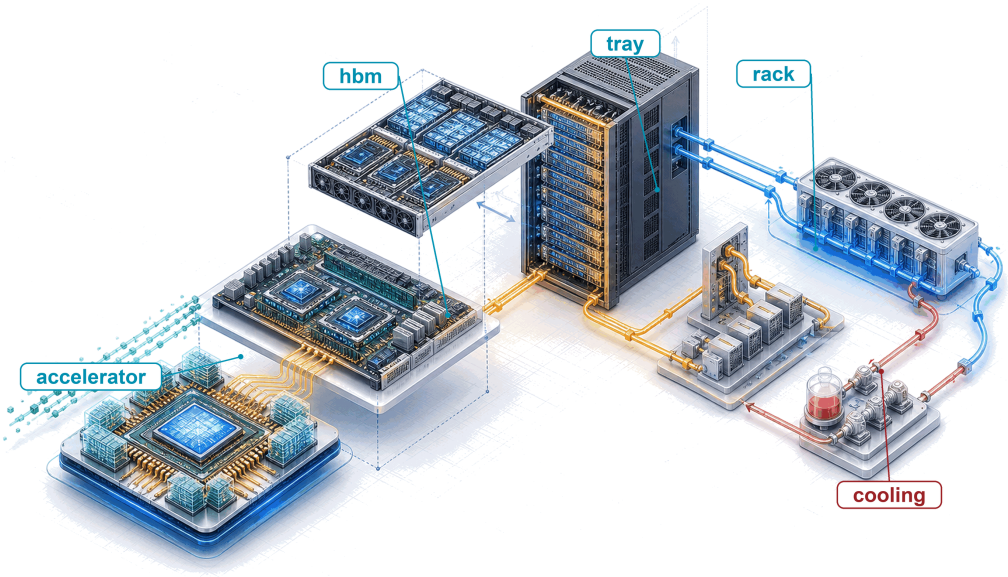

2.1 The Infrastructure Walls
2.2 The Accelerator
2.3 The Memory Wall
2.4 Sustained Compute and Precision
2.5 The Node
2.6 The Rack
2.7 The Pod
2.8 Accelerator Selection for Fleet Workloads
2.9 Infrastructure Technologies That Move the Walls
2.10 Fallacies and Pitfalls
2.11 Summary

Purpose

Why does infrastructure, not algorithms alone, determine who can participate in advancing machine learning at scale?

Machine learning systems have a physical reality beyond their code. Public algorithms do not make large-scale training equally accessible: executing them requires power, cooling, memory, and communication capacity that organizations must build and sustain. A model that outgrows one machine turns a local resource problem into a shared infrastructure commitment. Buying more accelerators adds arithmetic capacity, but it also adds data transfers, heat, and components whose failures can interrupt collective work. The useful capacity of the fleet depends on whether those demands can be met together. Hardware that appears attractive in isolation may spend much of its deployment waiting for data or operating below its advertised speed. A facility that can house equipment may lack the electrical or cooling capacity to run it at the intended density. These mismatches consume both capital and operating budgets while reducing the experiments that can finish on schedule. Infrastructure decisions also outlast individual models, so an organization must weigh the efficiency of specialization against the flexibility needed for future workloads. Single-machine performance remains part of that judgment, but multiplying it by the number of machines cannot establish what the complete installation will deliver. Participation at scale therefore depends on turning purchased hardware into sustained, usable capacity within physical and financial limits. In C^3 terms, the physical substrate sets the compute and communication bounds within which the fleet must coordinate useful work.

Learning Objectives

- Explain how the memory, communication, power, and reliability walls bind at successive physical scales and cascade upward
- Compare GPU, TPU, and ASIC specialization as a bet on workload stability against refresh cycles and facility lifetime
- Distinguish memory capacity from memory bandwidth as separate procurement tests, using roofline placement to find bandwidth-bound phases
- Calculate token-latency floors, memory budgets, and pod training time from accelerator specifications and sustained utilization
- Match tensor, pipeline, and data parallelism to the available interconnect bandwidth hierarchy
- Evaluate rack and pod designs using power density, cooling limits, topology, and facility reliability constraints
- Select accelerator infrastructure using sustained throughput, power constraints, workload fit, and deployment cost

2.1 The Infrastructure Walls

The scaling laws of section 1.3 are why demand outruns any single device. Loss falls predictably as parameters, training tokens, and compute grow together, so each step up in capability multiplies both the memory a model occupies and the arithmetic its training consumes; a GPT-3-class model already needs several accelerators merely to store its FP16 weights. A model that exceeds one accelerator's memory capacity forces a hardware decision before training begins. Adding accelerators supplies capacity, but also creates transfers between devices, increases power demand, and couples progress to more components that can fail. The familiar single-node compute and memory limits therefore become inputs to fleet design.

Four recurring physical constraints appear at escalating physical scales, from an eight-GPU node to a 25K-accelerator fleet (figure 2.1). Memory bandwidth, rack power, pod communication, and fleet failure exposure provide the examples; the workload and configuration determine which constraint binds at each scale.

At the scale moment, a model outgrows one machine and meets the memory, energy, and network walls, and growing the fleet opens its reliability gap. The four constraints here are those same limits, seen from the hardware that imposes them, so this chapter names each by the physical resource that runs out. The energy wall becomes the power wall, the limit on power delivery and heat removal. The network wall becomes the communication wall, the limit on link bandwidth. The reliability gap becomes the reliability wall, the same exposure measured as a failure interval. The memory wall keeps its name.

↩ The memory, network, and energy walls of the scale moment are established in section 1.1.

The 175B model anchors this progression. It is the chapter's running instance of Archetype A (GPT-4/Llama-3), a GPT-3-class dense transformer trained with mixed precision and the Adaptive Moment Estimation (Adam) optimizer, whose reference training cluster is 1,024 H100 GPUs in 128 eight-GPU nodes. Accelerator selection establishes the compute rate, memory capacity, and power envelope that a fleet can buy. The node must connect enough of those accelerators to hold and execute the workload; the rack must power and cool them; the pod must sustain communication and survive component failures. Each level's limit becomes an input to the next, because node power sets what the rack must cool and rack density shapes the pod's layout. This progression is the hardware contribution to C^3, the compute, communication, and coordination taxonomy of section 1.5. Useful fleet compute depends on the communication and coordination costs created at each physical boundary.

↩ The three fleet workload archetypes are established in section 1.6.1.

↪ The running model's parameters and reference cluster are developed in section G.1.

The fourth wall follows the naming convention of section 1.1, which separates two failure figures. A failed GPU is replaced rather than repaired, so the per-device figure is a mean time to failure (MTTF). Nodes and fleets keep running through repeated failures and replacements, so their figure is a mean time between failures (MTBF). The reliability calculation therefore starts from a per-GPU MTTF and ends at the MTBF of a node or fleet.

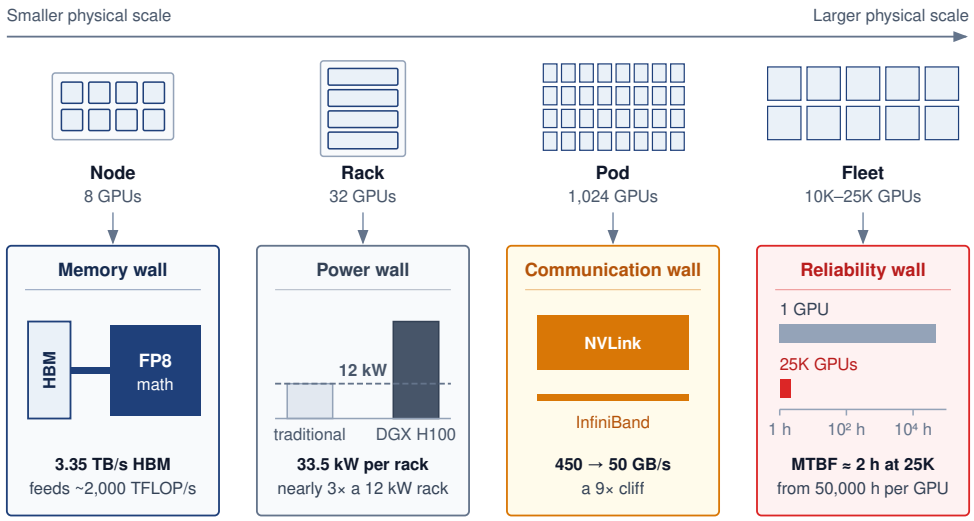

Figure 2.1: The Four Infrastructure Walls: Read left to right, the examples illustrate four constraints at increasing physical scales; the workload and configuration determine which one binds. The node compares 3.35 TB/s of HBM3 with about 2,000 TFLOP/s of FP8 arithmetic; the four-node DGX H100 rack draws 33.5 kW; the pod compares per-direction NVLink bandwidth with one InfiniBand port, a ninefold budget gap; and the fleet assumes independent GPU failures, reducing a 50,000-hour per-GPU mean time to failure (MTTF) to about 2 hours between failures at 25K GPUs.

The chapter climbs the same scale ladder, so the four walls arrive in the order the figure reads, though not one wall per section. The memory wall begins at the accelerator and is developed in section 2.3. The communication wall first appears inside the node, where section 2.5.1 measures the cliff between local and inter-node links, and returns at pod scale in section 2.7. The power wall binds at the rack (section 2.6.1), and the reliability wall accumulates from node health (section 2.5.5) through facility controls to the fleet-scale failure interval of section 2.7.1.

2.2 The Accelerator

The progression begins with the accelerator, because every wall above it is measured against the rates one device can compute, store, and draw. An accelerator purchase commits the fleet to a workload and software stack for longer than one experiment. A research group changing model operations frequently needs programmability; a service executing a stable computation at sustained volume can justify more specialization. The familiar trade-off between control flexibility and arithmetic efficiency matters here because an architectural mismatch is replicated across the fleet. Comparing programmable and specialized accelerators therefore starts with how much of the workload and compiler stack can remain stable over the deployment's lifetime.

Systems Perspective 2.1: The generality tax

General-purpose CPUs allocate silicon to caches and control logic as well as arithmetic. Branch predictors, reorder buffers, and instruction decoders support code with irregular control flow and data access. Matrix multiplication has a more regular schedule, so some of that flexibility contributes little to its dominant computation. The generality tax (principle 1) is the area and energy devoted to capabilities outside the dominant workload.

Specialized accelerators reduce this tax by replacing general-purpose control with a dataflow matched to the workload. In the approximate shares of figure 2.2, the arithmetic fraction of the die rises from roughly 5–10 percent on a CPU to 50–60 percent on a GPU, 70–80 percent on a Tensor Processing Unit (TPU), and over 90 percent on a custom application-specific integrated

circuit (ASIC). The shares mark a direction rather than measured proportions for any product. The selection question is how much of the complete workload can use the specialized hardware.

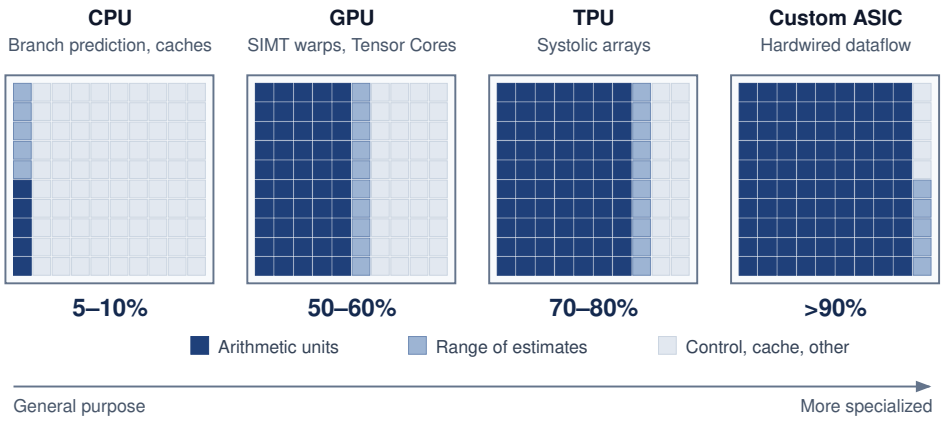

Figure 2.2: The Accelerator Spectrum: Each die is drawn as a grid of 100 cells divided between arithmetic units (dark) and control, cache, and other logic (light), with the medium cells marking the spread of each approximate, illustrative share. Read left to right, the arithmetic share grows from 5 to 10 percent on a CPU to over 90 percent on a custom ASIC, because hardwired dataflow replaces instruction fetch, decode, and scheduling; each step buys throughput per watt by narrowing the range of model architectures the silicon runs well.

Along the spectrum of figure 2.2, GPUs retain substantial programmability and provide 10–100× the matrix throughput of CPUs, while custom ASICs at the far end hardwire a specific dataflow for maximum efficiency at the cost of flexibility. The spectrum is not a ranking from "bad" (general-purpose CPU) to "good" (custom ASIC). It is a continuum of trade-offs in which each position is the best match for a different combination of workload characteristics and deployment constraints.

GPUs represent the first major step away from general-purpose computing toward large-scale parallelism while retaining programmability. An NVIDIA H100, for example, contains 16,896 CUDA cores organized into 132 streaming multiprocessors (SMs). The SMs keep thousands of threads in flight under the single instruction, multiple threads (SIMT) model, so a programmer-written kernel can express any mathematical operation, from convolutions to attention to custom loss functions. Neural network operations have uniform enough control flow that this programmability costs the GPU little on ML workloads.

TPUs specialize the matrix dataflow while retaining other units for operations that do not map onto it. Google's Tensor Processing Units use a **systolic array**[1] architecture, a fixed grid of multiply-accumulate (MAC) units that pass operands and partial results between neighboring cells on a regular schedule, so local reuse replaces per-operation instruction fetch and memory traffic. The complete accelerator still needs control, data movement, and execution of nonmatrix operations.

The cost of this efficiency is reduced programmability. Models must be compiled through Accelerated Linear Algebra (XLA), which maps high-level operations onto the fixed dataflow. Workloads with irregular control flow or dynamic shapes may not map well to this architecture, and the compilation process itself can be time-consuming (minutes to hours for complex models), which slows the iteration cycle during model development.

The programming model distinction between GPUs and TPUs has practical implications for organizational decisions. A research team that frequently adds custom operations, such as new attention patterns or activation functions, needs a path to implement them as kernels without waiting for compiler support, and must check how each platform expresses and optimizes those operations. A team running a stable computation at scale can place more weight on sustained utilization and compilation cost amortized over many runs. CUDA kernel support and XLA graph compilation offer different development paths, but neither establishes a utilization advantage without evaluating

[1] **Systolic array**: From Greek *systole* (contraction), the name describes data pulsing through a grid of processing elements in a heart-like rhythm. Kung and Leiserson (1979) proposed the design for VLSI chips whose limiting resource was memory access rather than logic (Kung 1982): each operand fetched once is reused by every cell it passes. That ratio of computation to memory traffic is exactly what matrix multiplication at ML scale rewards.

the workload and implementation. Software fit belongs in the hardware comparison because unsupported or poorly mapped work can leave specialized arithmetic idle.

Example 2.1: The TPU origin story

Scenario: In 2013, Google projected that if users spoke to Android phones for 3 minutes per day using voice search, handling the inference load on general-purpose CPUs would require doubling data center compute capacity at prohibitive cost (Jouppi et al. 2017).

Diagnosis: Standard CPU architectures paid a heavy generality tax, spending significant silicon area and power on out-of-order execution, branch prediction, and complex caches rather than matrix multiplication throughput.

Systems lesson: Hardware specialization is economically justified when a workload is large, stable, and dominated by a single mathematical kernel (MatMul). The TPU v1 hardwired a 256×256 systolic array to deliver 92 TOPS (INT8) within a 75 W footprint, establishing custom ASIC accelerators in production serving infrastructure.

Custom ASICs occupy the specialized end of the spectrum, where the economics of silicon can justify giving up much of the general-purpose programmability. When an organization runs a stable workload at enormous scale, designing a chip specifically for that workload may become economically rational. A custom ASIC can improve energy per operation and sustained utilization on its target workload by removing features that computation does not require.

Tesla's Dojo D1 chip, for example, is optimized for the spatial and temporal structure of video-based vision models. Its hundreds of tightly coupled processing nodes are designed around the dataflow needs of Tesla's vision pipeline, with on-chip SRAM sized to keep spatial tiles close to the compute units. This reduces the repeated round-trips to off-chip memory that a general-purpose GPU would require for the same computation. AWS Trainium takes a different approach, targeting the broad category of Transformer training rather than a single model. Trainium systems combine compiler-managed memory scheduling with hardware-assisted collective communication, so common training patterns such as tensor-parallel synchronization and data-parallel reductions can be optimized in the accelerator fabric rather than handled entirely by host software.

The risk of custom silicon follows from the same specialization. A change in the dominant model architecture can reduce the value of a fixed dataflow if the new operators, tensor shapes, or memory demands map poorly onto it. Sharing matrix operations across model families can preserve some of that investment, so a change from convolutional neural networks (CNNs) to Transformers does not by itself make an accelerator unusable. The forecasting window is long. Even Google's first TPU, a high-priority project, took 15 months to design, verify, build, and deploy (Jouppi et al. 2017), and a chip then serves for several years more inside a facility planned to outlast multiple accelerator generations (section 2.9). The design decision must therefore anticipate which operations will remain useful years after the workload was profiled, and how much flexibility is needed for the rest. In 2015, few would have predicted that attention-based Transformers would replace CNNs as the dominant architecture within five years, a span of only a few ASIC design cycles.

Custom silicon is therefore a bet on workload stability, and the organizations that make this bet are typically those with enough scale to justify a large up-front design investment and enough workload volume to amortize that NRE (nonrecurring engineering) cost across many chips. Google can justify the TPU's development cost because high-volume services such as Search, YouTube recommendations, and Gmail spam filtering can amortize a shared accelerator platform. A research lab running a few hundred GPUs cannot justify the same investment.

The bet can pay off when workloads are stable. The first-generation TPU reported 30–80$\times$ the performance per watt of contemporary GPUs and CPUs on its inference workloads (Jouppi et al. 2017), and TPU v4 reported 1.3–1.9$\times$ lower power than the A100 for similar-sized training systems (Jouppi et al. 2023). A workload shift can strand that advantage because the chip's fixed dataflow may not accommodate a substantially different computational pattern. A different specialization bets on physical scale rather than a fixed dataflow, expanding the compute fabric to the size of an entire wafer.

2.2.1 Wafer-scale engines

Slicing a silicon wafer into discrete chips forces every exchange between accelerators to leave the die and cross slower, higher-energy package, board, or rack links. Wafer-scale engines (WSE) avoid those crossings by keeping most of a silicon wafer intact as one continuous compute fabric the size of a dinner plate, as in the Cerebras WSE-3, rather than dicing the wafer and reconnecting the pieces across packages and boards.

For its 2024 WSE-3 configuration, Cerebras reported 900,000 AI-optimized cores, 44 GB of distributed on-chip SRAM, and 21 PB/s of aggregate memory bandwidth in its registration filing (Cerebras Systems 2024). These specifications are not directly comparable to GPU-cluster peak FLOP/s or high-bandwidth memory (HBM) bandwidth because the architectures expose different dataflows, memory organizations, precisions, and workload mappings. The durable comparison is topological. On-wafer links avoid package-, board-, and rack-level transfers for work that maps onto the fabric.

A multi-GPU chassis and a wafer-scale engine draw their package and board boundaries in different places (figure 2.3). The eight GPUs of a chassis exchange state over NVLink, NVIDIA's GPU-to-GPU link, and NVSwitch chips join those links into one fabric. Transfers to other nodes cross the external network. A wafer-scale engine instead routes between 900,000 on-wafer cores on the 300 mm wafer. Defect-aware routing lets the fabric use working regions despite local manufacturing flaws. In the documented CS-2 configuration, twelve power supplies in a 9+3 redundant arrangement could deliver up to 23 kW to the WSE-2 (Systems 2021). That configuration-specific design illustrates why wafer-scale power delivery and cooling require system-level liquid-cooling hardware rather than a conventional accelerator heatsink.

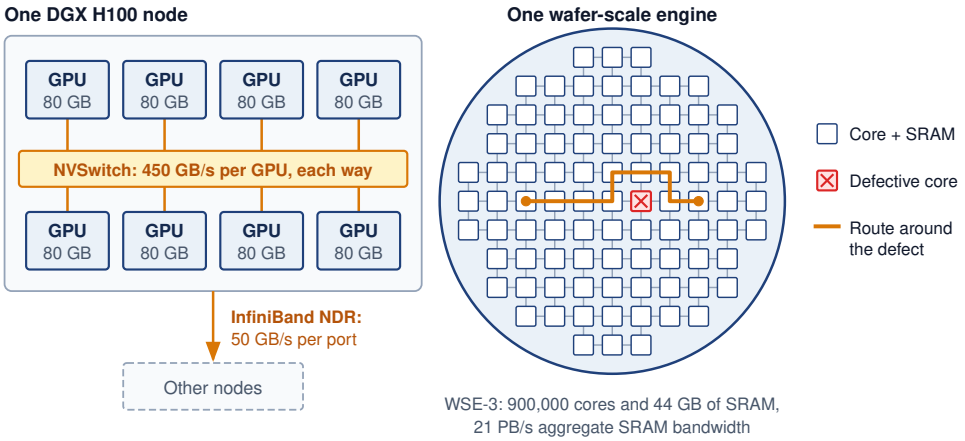

Figure 2.3: Wafer-Scale Engine (WSE) Architecture: A DGX H100 node (left) joins eight separately packaged GPUs through an NVSwitch fabric and must leave the chassis over InfiniBand to reach other nodes, so every exchange crosses a package, board, or network boundary. A WSE (right) instead uses most of a 300 mm wafer as one routed mesh of cores, each with local SRAM; defect-aware routing steers traffic around unusable cores, and the on-wafer mesh replaces off-package transfers for mapped operations. The WSE-3 figures are the vendor's reported totals and are not directly comparable to GPU HBM bandwidth.

Wafer-scale engines specialize the *physical* architecture while retaining flexibility in the *computational* model through programmable cores. That trade-off exchanges manufacturing complexity and defect-aware routing for an on-wafer fabric that replaces package- and board-level communication between its cores, though SRAM capacity and data injection still constrain which parts of a workload can benefit from that locality. Where the wafer-scale engine removes the package boundary by never dicing the wafer, chiplet designs keep the package and change what goes inside it.

2.2.2 Chiplet-based accelerators

Once a monolithic processor die grows to the optical exposure limit of a lithography scanner, adding more arithmetic units or HBM controllers on a single piece of silicon becomes impossible without

stitching exposures at lower yield and sharply higher cost. Chiplet architectures, exemplified by NVIDIA's Blackwell and AMD's Instinct MI300 series, break through that ceiling by partitioning the processor into smaller dies connected by a high-bandwidth die-to-die interconnect on a common package substrate. This approach addresses two physical limits that constrain monolithic designs.

First, the maximum die size is capped by the photolithographic reticle limit, the largest area that a lithography scanner can pattern in a single exposure. The limit comes from the optics of photolithography, because the lens system of an extreme-ultraviolet (EUV) scanner projects a pattern onto a field of only about 26×33 mm (858 mm^2). A larger die would need multiple exposures stitched together, which is technically possible but sharply raises cost and lowers yield. The limit caps the number of Tensor Cores (the GPU's matrix-multiply units), SMs, and HBM stacks a monolithic GPU can integrate. Chiplet designs bypass this limit by placing multiple dies on a single package, with the B200's two reticle-limited dies effectively creating a processor of about 1,700 mm^2.

Second, manufacturing yield depends on die area and defect density. Under the standard Poisson defect model, the fraction of defect-free dies falls exponentially with die area at a fixed defect density, because a single killer defect anywhere on the die can render it unusable. Smaller dies expose less area to each defect, and testing can select known-good chiplets before package assembly, so one defect discards a small die rather than a reticle-sized one. The resulting yield advantage must be weighed against the cost and yield of the package, its interconnects, and the additional testing. Chiplets change the manufacturing trade-off; they do not guarantee a lower cost per unit of compute.

The trade-off is the die-to-die interconnect. Communication between chiplets within a package is faster than communication between packages (NVLink) but slower than communication within a single monolithic die (the on-die mesh network). Workloads that generate frequent, fine-grained communication between processing elements (such as operations that share data between nonadjacent SMs) may experience a latency penalty when those SMs reside on different chiplets. GPU vendors mitigate this by making the die-to-die link transparent to the programmer, so the software sees a single logical GPU regardless of the underlying chiplet topology. Seeing why modern accelerators reached that multi-die boundary requires tracing how successive GPU and TPU generations spent their single-die transistor budgets first.

2.2.3 Evolution of GPU architectures

NVIDIA's data center GPUs show the specialization bet placed one generation at a time. Each generation added hardware aimed at a bottleneck that production ML workloads had exposed, rather than more transistors of the same kind. The Volta architecture (2017) introduced the first-generation Tensor Core, dedicated matrix-multiply-accumulate hardware placed beside the existing CUDA cores, because neural network training spends most of its time in matrix multiplications. Volta accelerated that dominant pattern without giving up the general-purpose programmability that made GPUs attractive for research.

The Ampere architecture (2020) expanded Tensor Core support to TF32 (a format with FP32's range and less precision) and BF16 and introduced Multi-Instance GPU (MIG), which partitions one GPU into up to seven isolated instances so that several inference workloads can share it. The change that mattered most for infrastructure was Ampere's third-generation NVLink, which doubled per-GPU bandwidth to 600 GB/s from Volta's 300 GB/s, the bandwidth that tensor parallelism needs because it splits each layer's matrices across GPUs.

The Hopper architecture (2022) added the Transformer Engine, whose FP8 recipes and autocast integration, the software that selects each operation's precision, run supported Transformer operations in FP8 with managed scaling while retaining higher precision where the configured recipe or implementation requires it. Hopper also introduced NVLink 4.0 at 900 GB/s per GPU and the NVLink Switch, enabling NVLink connectivity beyond the 8-GPU boundary.

The Blackwell architecture (2024) takes the chiplet route of section 2.2.2 past the 858 mm^2 reticle limit. The B200 pairs two GPU dies in one package over a 10 TB/s die-to-die link, fast enough that the two dies appear to software as a single logical GPU for many operations, and its fifth-generation NVLink doubles intra-node bandwidth again to 1800 GB/s per GPU. Its rating in table 2.1 is an FP8 figure, so a comparison at FP16 or BF16 should use the vendor's lower rate for that precision.

Each generation thus targeted the bottleneck its predecessor exposed. Volta added dedicated matrix hardware, Ampere widened mixed-precision support, Hopper added FP8 for Transformer

workloads, and Blackwell worked around the die-size limit with two dies in one package. Because architects choose those targets by profiling the workloads of their era, accelerator design follows model architecture, which is why Transformer workloads have shaped accelerator design since 2017.

Because accelerator generations turn over faster than the facility that houses them, a data center must anticipate refresh cycles before the first rack arrives. Those generations also changed interconnect capacity and power (table 2.1). From V100 to B200, dense FP16 throughput rose 18×, memory bandwidth 8.9×, and thermal design power (TDP) 3.3× (figure 2.4). These are growth multiples over the whole span, not annual rates, and their ratios carry the planning lesson. Each byte delivered from memory must now feed about twice as many operations, and each watt of the power allowance supplies more than five times as many peak operations. That first ratio raises the data reuse a workload needs to approach peak throughput. Even with the fivefold gain in peak operations per watt, reaching the 18× compute multiple still required raising per-device TDP by 3.3×, and that higher per-device wattage multiplied by device count sets the rack load. A facility that replaces accelerators one for one to gain compute density therefore outruns both its memory delivery and its cooling envelope.

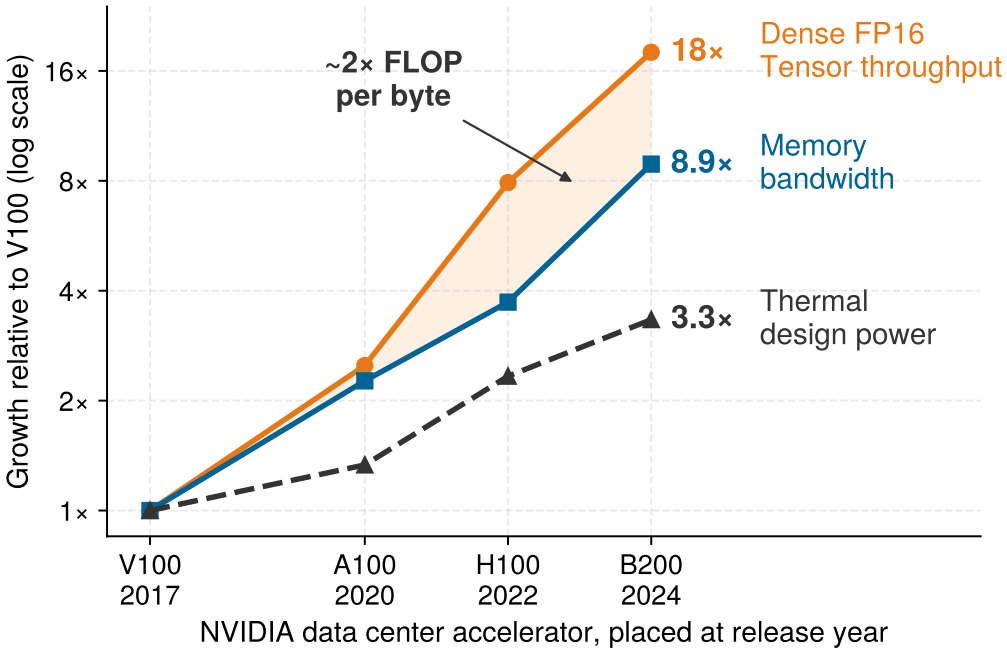

Figure 2.4: Compute, Bandwidth, and Power Scale at Three Different Rates: Dense FP16 Tensor throughput, memory bandwidth, and thermal design power for NVIDIA data center accelerators, normalized to the V100 on a log scale, where the shaded gap between the compute and bandwidth lines is the growth in arithmetic intensity a kernel needs to stay compute-bound. Compute rises 18×, bandwidth 8.9×, and TDP 3.3×, so the endpoint gains in arithmetic exceed the gains in memory bandwidth and power allowance. These are peak specifications, not measured application throughput. Specifications come from the MLSysIM hardware registry.

Table 2.1: NVIDIA GPU Architecture Evolution: Each generation introduces architectural features targeted at the dominant ML workload pattern of its era. Efficiency rises substantially but unevenly across generations, while NVLink bandwidth increases from 300 GB/s on V100 to 1,800 GB/s on B200, with a smaller 1.5× A100-to-H100 step between the larger doubling jumps.

Architecture	Year	Key innovation	Peak tensor	TDP	NVLink BW
Volta (V100)	2017	First Tensor Core	125 TFLOP/s FP16	300 W	300 GB/s

Architecture	Year	Key innovation	Peak tensor	TDP	NVLink BW
Ampere (A100)	2020	Multi-precision, MIG	312 TFLOP/s FP16	400 W	600 GB/s
Hopper (H100)	2022	Transformer Engine, FP8	1,979 TFLOP/s FP8	700 W	900 GB/s
Blackwell (B200)	2024	Dual-die, NVLink 5	4,500 TFLOP/s FP8	1000 W	1,800 GB/s

Because table 2.1 mixes FP16 and FP8 peaks, its per-watt efficiency must be read at fixed precision. At dense FP16, dividing each generation's peak throughput by its TDP (table 2.2) gives a V100-to-B200 gain of 5.4× in peak throughput per watt, the more-than-fivefold figure above, against 18× in raw throughput. The rest of the raw gain came from a power allowance that itself rose 3.3×. The same caution applies between adjacent generations. The H100's FP8 peak per watt of TDP is 6.8× the V100's FP16 figure, but a factor of two of that comes from the change of precision, and at equal FP16 precision the gain is 3.4×. Facilities must therefore accommodate higher device power as well as greater arithmetic capacity, which connects this trend to the power wall in section 2.6.1.

Table 2.2: Power Efficiency across GPU Generations: Dense FP16 tensor arithmetic peak divided by accelerator TDP, with no sparsity multiplier. These ratios compare arithmetic ceilings, not measured training throughput or facility energy. Matched workload measurements must also include achieved utilization, memory and communication stalls, and non-accelerator power.

Generation	Dense FP16 peak TFLOP/s	TDP (W)	TFLOP/s per watt	Relative efficiency
V100 (2017)	125 TFLOP/s (FP16)	300 W	0.42 TFLOP/s/W	1×
A100 (2020)	312 TFLOP/s (FP16)	400 W	0.78 TFLOP/s/W	1.9×
H100 (2022)	989 TFLOP/s (FP16)	700 W	1.41 TFLOP/s/W	3.4×
B200 (2024)	2250 TFLOP/s (FP16)	1000 W	2.25 TFLOP/s/W	5.4×

The four generations in table 2.2 sit on a broader and slower trend. Across the 64 accelerator models released since 2017 for which Epoch AI records both a dense FP16 tensor rate and a TDP, from every major vendor, peak arithmetic per watt improved about 1.2× per year, doubling roughly every 4 years (figure 2.5) (Epoch AI 2026b). The leading NVIDIA parts improved faster, 5.4× from the V100 to the B200 in seven years (the fixed-precision gain noted earlier), because frontier devices adopt new process nodes and tensor formats first. A fleet on a 4-year refresh cycle therefore gains about 2× in peak arithmetic per watt each cycle on the broad trend. Peak ratios bound what a refresh can deliver, and only measurements at the workload level decide how much of it a given fleet realizes.

↪ How these peak ratios screen a fleet refresh decision is developed in section 12.9.3.1.

2.2.4 Evolution of TPU architectures

Faster per-chip arithmetic leaves a cluster stalled if inter-chip collectives still have to traverse host CPUs and general-purpose PCIe switches. Google's TPU trajectory attacked that bottleneck directly by treating the multi-rack pod, rather than the individual chip, as its unit of design from its second generation onward. While GPUs emphasize peak per-chip throughput and software flexibility (CUDA), every TPU chip since TPU v2 integrates dedicated inter-chip interconnect (ICI) links that route traffic directly between accelerators without touching the host CPU.

The generation sequence shows how that pod-first design moved from inference to training. After the inference-only TPU v1 of the origin story, TPU v2 (2017) and TPU v3 (2018) shifted the architecture toward training. They added bfloat16 (BF16), which avoids the dynamic range problems of standard FP16, and they scaled the pod concept until a TPU v3 pod joined 1,024 chips with about

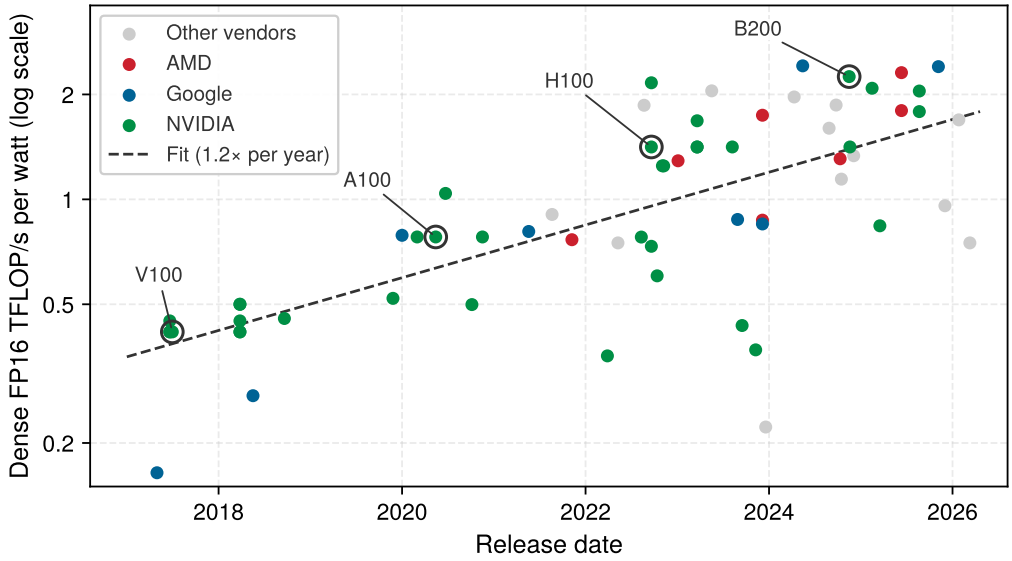

Figure 2.5: Peak Arithmetic per Watt over Time: Dense FP16 tensor throughput divided by TDP for ML accelerators released since 2017, on a logarithmic axis, colored by vendor. The dashed line is a least-squares fit across all devices; the four GPUs of table 2.2 are circled and labeled. The ratio compares arithmetic ceilings, not measured energy per workload. Data from Epoch AI (2026b).

126 PFLOP/s of aggregate BF16 peak. TPU v4 (2021) extended the same logic into the network itself. It added an optical circuit switch so the physical topology could be reconfigured for different workloads, introduced SparseCore units that accelerate embedding lookups, and raised per-chip BF16 throughput to 275 TFLOP/s.

The TPU v5p (2023) continued this pod-scale design for high-end training. It features 459 TFLOP/s of BF16 compute, 95 GB of HBM2e with 2,760 GB/s memory bandwidth, and 1,200 GB/s bidirectional ICI bandwidth per chip. The ICI figure matters for hardware selection because collectives such as AllReduce run over those links. AllReduce combines values such as gradient sums across a group of accelerators and returns the result to every participant. Section D.2.1 develops the complexity model and algorithm choices that determine when those exchanges limit training throughput. The pod-first design took shape over successive TPU generations (table 2.3).

Table 2.3: Google TPU Architecture Evolution: The TPU's development highlights a "pod-first" design philosophy where inter-chip interconnect bandwidth and compiler-level optimization (XLA) are as important as peak per-chip arithmetic. Note: TPU v1 was INT8 only.

Generation	Year	Key innovation	Peak BF16	HBM	ICI BW (aggregate, one direction)
TPU v1	2015	Systolic array (inference)	92 TOPS*	—	—
TPU v2	2017	High-bandwidth memory	46 TFLOP/s	16 GB	248 GB/s
TPU v3	2018	Liquid cooling, pod scale	123 TFLOP/s	32 GB	280 GB/s
TPU v4	2021	Optical switch, SparseCore	275 TFLOP/s	32 GB	300 GB/s
TPU v5p	2023	Larger pods, HBM2e	459 TFLOP/s	95 GB	600 GB/s

The GPU and TPU lines follow two design philosophies (table 2.1, table 2.3). GPU evolution packs more arithmetic, in ever lower-precision Tensor Core formats, onto one package, and turns to

chiplets once a single die reaches its size limit. TPU evolution treats the pod as the product. Per-chip arithmetic matters, but inter-chip bandwidth and reconfigurability decide how much of it a large training job can use. Both lines combine local computation with memory and interconnect improvements; they place those resources differently. The warehouse-scale perspective of section 1.2.2 treats chip throughput, collective bandwidth, and reconfigurability as parts of one system (Barroso et al. 2019), so the useful comparison is how each candidate sustains the intended workload across that complete system.

2.2.5 Matrix units

Both lines concentrate their peak arithmetic in specialized matrix units, called Tensor Cores on NVIDIA GPUs and Matrix Multiply Units (MXUs) on Google TPUs. A Tensor Core executes matrix-multiply-accumulate instructions of the form $\mathbf{D} = \mathbf{A} \times \mathbf{B} + \mathbf{C}$ on small tile matrices, concentrating the dominant neural-network operation in dedicated hardware that supplies most of the chip's peak arithmetic.

With 528 Tensor Cores across the chip (4 per SM, 132 SMs), the H100 reaches its 989 TFLOP/s dense FP16/BF16 peak. The important systems point is the architectural specialization rather than the exact per-cycle instruction accounting. The die dedicates a large fraction of its arithmetic area to dense tiled matrix operations rather than general-purpose scalar execution.

Google's MXUs organize multiply-accumulate units into the systolic array introduced with the TPU above. In the MXU's weight-stationary dataflow, a weight tile stays resident in the array while activations stream through it. Every operand fetched from memory is therefore reused by an entire row or column of cells, which is the reuse that a Tensor Core kernel obtains by staging tiles in shared memory and registers. Each TPU v5p chip contains two TPU cores (which Google calls TensorCores, a name unrelated to NVIDIA's Tensor Cores) with four 128×128 MXUs each, together providing 459 TFLOP/s of BF16 throughput. Approaching that peak requires compatible tile shapes and a sustained operand supply; pipeline fill, partial tiles, and dependencies can leave cells idle. Vector and scalar units handle other operations, so complete-workload efficiency depends on how computation maps across the accelerator.

For our 175B model, compiler and kernel support determine how much work reaches the specialized arithmetic units. CUDA kernels can combine matrix operations with thread-level work, while XLA maps the computation graph onto the TPU dataflow. In either stack, fusion keeps intermediate results in registers or shared memory instead of repeatedly writing them to HBM. Sustained throughput therefore depends on operation coverage and data reuse as well as the matrix units themselves.

Definition 2.1: Tensor core

Tensor Core is a specialized ML accelerator hardware unit that performs a fused matrix-multiply-accumulate operation $\mathbf{D} = \mathbf{A} \times \mathbf{B} + \mathbf{C}$ on small tiles (for example, $16 \times 8 \times 16$ in BF16) as a single hardware instruction, delivering dramatically higher throughput than general-purpose CUDA cores by trading programmability for fixed-function matrix arithmetic.

1. **Significance**: Tensor Cores provide the bulk of the H100's 989 TFLOP/s FP16/BF16 peak (and ~1,979 TFLOP/s at FP8), roughly 14.8× the ~67 TFLOP/s delivered by CUDA (vector) cores alone. This comparison changes both the execution unit and precision. Reductions and element-wise operations in layer normalization or softmax require a different performance model; their throughput is not a fixed fraction of the matrix peak.
2. **Distinction**: A Tensor Core is a fixed-function matrix-multiply-accumulate unit, not a faster general-purpose core. CUDA cores execute arbitrary scalar and vector work; Tensor Cores execute only the tiled $\mathbf{D} = \mathbf{A} \times \mathbf{B} + \mathbf{C}$ primitive, which is why peak throughput and delivered throughput diverge so sharply by workload.
3. **Common pitfall**: A frequent misconception is that all GPU operations benefit from Tensor Cores. Operations must map to the $\mathbf{D} = \mathbf{A} \times \mathbf{B} + \mathbf{C}$ tile computation to use those units. Element-wise activations, normalization, and embedding lookups have different

execution and data-access demands. A fixed 10–30× penalty per FLOP is not a workload-independent model for their runtime.

Table 2.4: The Accelerator Spectrum: The architectures place arithmetic, control, and memory differently. Specialization can improve throughput per watt when the workload and software stack use its dataflow; the table does not rank measured efficiency. Wafer-scale engines replace package and board boundaries with an on-wafer fabric of programmable cores.

Feature	CPU	GPU (H100)	TPU (v5p)	Wafer-scale	Custom ASIC
Arithmetic core	Scalar/vector	Tensor Core	Systolic array	RISC-style core	Fixed dataflow
Execution	Instruction	SIMT	Data-driven	Dataflow	Hardwired
Memory control	Cache	L1/L2 + HBM	Scratchpad	On-chip SRAM	Compiler-managed
Flexibility	Extreme	High	Moderate	High	Low

Across the whole spectrum, from CPU to custom ASIC, the architectures differ in arithmetic core, execution model, and memory control (table 2.4). For the 175B running model, the choice among these columns is not purely about peak FLOP/s. A research lab experimenting with novel architectures weekly can justify the GPU's generality tax by its flexibility. An organization deploying a fixed Transformer at scale for years may find that the TPU's dataflow efficiency or a custom ASIC's power advantage dominates total cost. The accelerator spectrum is ultimately an economic question of *how much flexibility can be surrendered, given the stability of the workload*. Even when the H100's arithmetic engine reaches 1,979 TFLOP/s in its advertised low-precision mode, eligible operations approach that throughput only when data reaches the compute units fast enough. Accelerator selection therefore requires a second comparison, between the workload's data demand and the memory system that feeds the arithmetic.

2.3 The Memory Wall

Single-request inference on our 175B running model has a latency floor set by memory delivery. Its weights occupy 350 GB in FP16. In dense autoregressive decode at batch size 1, each generated token uses all of those weights, whose footprint far exceeds on-chip storage, so every token must stream them in from memory. Suppose, hypothetically, that one H100 could hold them. Streaming the whole 350 GB through that device's HBM at about 3.35 TB/s takes over 100 ms per token, long enough to be a perceptible pause between generated words. No real H100 holds the model, so the figure is a floor on bandwidth demand rather than a prediction of deployed latency, but it shows that the memory wall is a delay a user can notice, not an abstract architectural ratio.

That floor reflects a structural divergence. Transistor scaling has multiplied logic throughput far faster than off-chip interconnects can supply operands. In phases with little data reuse, where each fetched byte feeds few operations, specialized matrix units therefore routinely spend more time waiting for bytes than executing math. This divergence is the memory wall (Wulf and McKee 1995). In those phases the weight and activation streams of large neural networks make memory bandwidth rather than peak compute the governing constraint.

↪ The diagnostic table that maps memory-bound, compute-bound, and latency-bound symptoms to the optimizations that help is developed in section A.3.2.

Three levers ease that constraint. Sharding spreads the weights across more memory paths, so each path streams fewer bytes. Quantization shrinks the bytes themselves. Reuse lets each fetched byte feed more arithmetic, which batching and on-chip tiling make possible. The Roofline Model, recapped for one machine in section A.3.1, tests whether the resulting data supply can feed the arithmetic units. Its ridge point, peak arithmetic divided by memory bandwidth, is the number of operations a kernel must perform per byte fetched before arithmetic rather than memory becomes the limit. The ridge point recurs throughout the analysis below, and section 2.3.2 evaluates it for the H100. Where each architecture places its bytes sets how much bandwidth those levers can draw on (figure 2.6), because moving the active working set closer to arithmetic increases available bandwidth while reducing the capacity that can reside there.

Accelerators do not line up along a single performance axis; they fall into structural clusters according to where each architecture places its bytes (figure 2.6). Edge devices in the lower band

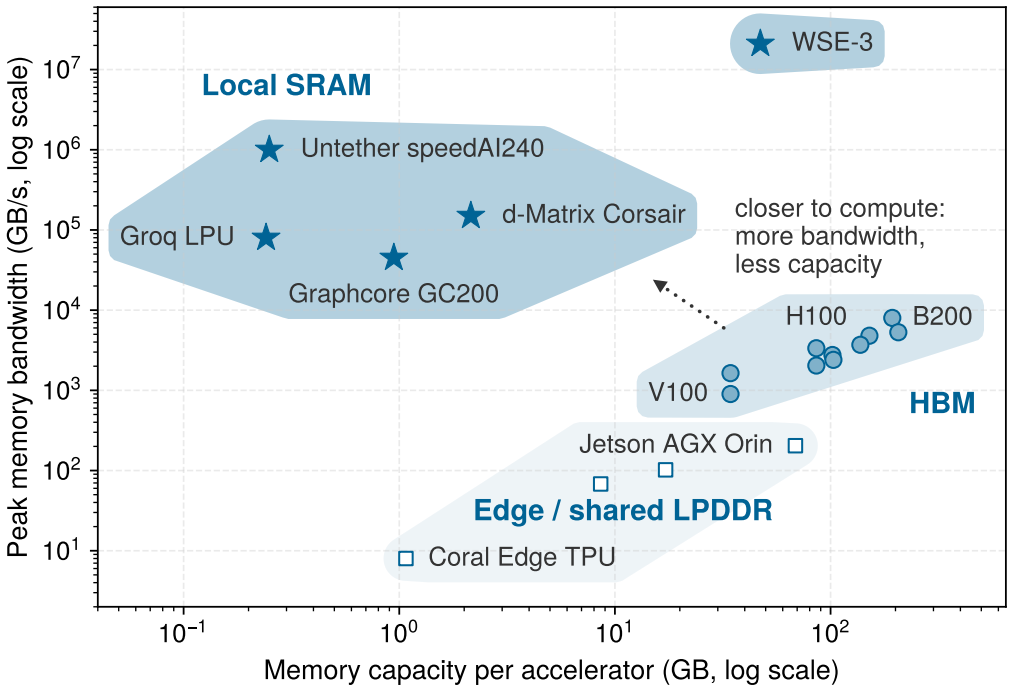

Figure 2.6: Accelerator Memory Tiers: Peak memory bandwidth and capacity separate local-SRAM designs (solid stars), HBM accelerators (half-tone filled circles), and edge and shared-LPDDR devices (hollow squares); all three are memory, so they share one hue and differ by marker. The dotted arrow marks the design trade-off: moving data closer to compute raises bandwidth but reduces capacity.

keep memory cheap and compact, but their bandwidth remains constrained, and the devices of Archetype C (Federated MobileNet) sit in that corner. The binding limit on each such device, however, is not its place on this memory map but the power envelope that bounds compute. A phone-class device sustains only 2 W to 3 W, while one H100's thermal envelope is about 233× larger. A federated fleet therefore grows by adding devices rather than by moving up the memory tiers, and at fleet scale, as section 1.6.1 noted, communication and which devices participate in each round limit progress as well.

Within the data center tiers, HBM moves memory onto the accelerator package, buying enough capacity and enough bandwidth for active model state. Local-SRAM designs push the working set still closer to arithmetic, buying extraordinary bandwidth at the cost of much smaller resident capacity. Of the two, HBM is the next engineering response, because it is the highest-bandwidth tier that can hold tens to hundreds of gigabytes next to the compute die, although it is not the fastest memory in the system.

2.3.1 HBM and the memory wall

An H100 can execute 1,979 TFLOP/s of low-precision matrix arithmetic, but the batch-1 weight-streaming example for our 175B model supplies too little reuse to approach that peak. In a bandwidth-floor model its arithmetic units would use well under 1 percent of the matrix peak while the weights stream in. The limiting resource is the speed of *delivery* relative to *multiplication*. Once weights fit through sharding or compression, the serving path must still deliver the active shards at the required rate. Additional arithmetic capacity alone does not raise a binding bandwidth ceiling. HBM increases that data supply within the package.[2]

HBM is the bottom on-package tier of the accelerator memory hierarchy, beneath the register file, the L1 and shared memory (SRAM),[3] and the L2 cache. Aggregate bandwidth falls from hundreds of TB/s at the registers to the 3.35 TB/s of the H100's 80 GB of HBM3, while capacity grows in the opposite direction.

↪ The power, thermal, and memory limits of phone-class edge devices are developed in section 11.3.

2 **HBM origin**: Standardized by JEDEC in 2013 as a joint development between AMD and SK Hynix, originally for graphics cards. ML accelerators adopted HBM because neural networks share the same bandwidth-hungry, capacity-moderate access pattern as high-end rendering. Each HBM generation has roughly doubled bandwidth (128 GB/s in HBM1 to 1.2 TB/s per stack in HBM3e), yet the gap between memory bandwidth and arithmetic throughput continues to widen, which makes HBM a necessary but never sufficient response to the memory wall.

The bandwidth gap between registers and HBM is approximately 100×, but that ratio alone does not specify stalled time. Utilization depends on **arithmetic intensity**, the floating-point operations performed per byte moved from memory (FLOP/byte), which tiling raises by staging matrix blocks on chip so that each fetched operand feeds many multiply-accumulate operations. Section A.5 collects tiling with the other kernel-level levers for reference.

Definition 2.2: HBM

HBM is a 3D-stacked DRAM architecture used by ML accelerators, with through-silicon vias (TSVs) connecting memory dies vertically within each stack and a wide interface across the package interposer connecting the stack to the processor.

1. **Significance**: HBM3 in the H100 delivers approximately 3.35 TB/s of memory bandwidth (roughly 17× DDR5's ~200 GB/s) at roughly 4 pJ/bit for an HBM2-class stack (O'Connor et al. 2017) vs. roughly 20 pJ/bit for off-package DRAM (Horowitz 2014). This bandwidth sets the BW term of the single-machine iron law given in section A.2. The H100's FP16 ridge point is approximately $989\,\text{TFLOP/s} / 3.35\,\text{TB/s} \approx 295$ FLOP/byte, so any operator with arithmetic intensity below 295 FLOP/byte (for example, single-request decode, which performs two FLOPs for every two-byte FP16 weight it reads) remains memory-bound even with HBM.
2. **Distinction**: Unlike socketed DDR memory, which communicates across centimeters of printed circuit board (PCB) trace, an HBM stack sits beside the processor on the package. Its micrometer-scale vertical connections can be 1,000× shorter than board traces, but the complete processor-to-memory path also crosses the interposer. The wide, short on-package interface raises bandwidth without relying on a proportionate increase in per-pin signaling rate.
3. **Common pitfall**: A frequent misconception is that HBM solves the memory wall. HBM moves the wall rather than eliminating it. From the V100 (900 GB/s with 125 TFLOP/s FP16) to the H100 (3.35 TB/s with 989 TFLOP/s FP16), peak arithmetic (R_{peak}) grew faster than HBM bandwidth.

Traditional DDR memory connects to the processor through pins on the edge of a printed circuit board. Each DIMM communicates over a 64-bit bus, and even with multiple channels (8 channels is typical for a high-end server CPU), a modern server tops out at roughly 200 GB/s of aggregate memory bandwidth.

The physical distance from the DIMM slot to the processor die is measured in centimeters, and every centimeter of copper trace introduces capacitance, signal attenuation, and energy loss. At DDR5 data rates (4,800–6,400 MT/s per pin), the signal conditioning circuits must compensate for that degradation, consuming substantial power per bit transferred. Increasing the data rate on these long traces requires progressively more power for signal conditioning, creating a diminishing-returns curve that DDR5 is already approaching.

HBM sidesteps that board-level degradation through 2.5D integration (figure 2.7). Instead of driving signals horizontally across centimeters of motherboard trace over a narrow 64-bit bus, it stacks DRAM dies vertically and places the stack beside the processor on a silicon interposer.

The vertical stacking changes how bandwidth scales. Rather than pushing signals faster through long copper traces (the DDR approach, which has diminishing returns), HBM multiplies the number of parallel signal paths. A single HBM stack bonds 8–12 DRAM dies, threading thousands of through-silicon vias (TSVs),[4] copper pillars etched through each die, to form a 1024-bit-wide interface per stack, compared to 64 bits for a DDR5 channel. This 16× wider interface is the main source of HBM's bandwidth advantage; HBM need not exceed DDR's per-pin signaling rate. The processor and HBM stacks share a package and communicate across the silicon interposer, shortening the electrical path relative to socketed DIMMs. This packaging primarily improves bandwidth and energy per transferred bit; it does not by itself imply that complete HBM load latency is only a few nanoseconds.

The shorter path also lowers transfer energy per bit by roughly 5×, from approximately 20 pJ/bit for off-package DRAM (Horowitz 2014) to about 4 pJ/bit for HBM2 (O'Connor et al. 2017), because

[3] **SRAM energy (pJ/bit)**: Reading a bit from an 8–32 KB on-chip SRAM array costs roughly 0.15–0.3 pJ at the 45 nm node (Horowitz 2014), while fetching a bit from HBM2 costs about 4 pJ (O'Connor et al. 2017). This order-of-magnitude energy difference is the physical driver behind kernel fusion: every intermediate result kept in SRAM instead of written to HBM saves both time and significant power and thermal headroom.

↪ The single-accelerator memory hierarchy, from registers and shared memory through HBM to host memory, is collected for reference in section A.4.

↪ Arithmetic intensity and the boundary it draws between memory-bound and compute-bound work are collected for reference in section A.3.

[4] **TSV (through-silicon via)**: A vertical copper pillar, 5–10 micrometers in diameter, etched through a silicon die to connect stacked layers. Originally developed for CMOS image sensors in smartphone cameras, TSVs enabled HBM by replacing centimeters of PCB trace with tens of micrometers of vertical silicon, roughly a 1,000× shorter signal path, which is what lets a stack expose a 1024-bit interface within an accelerator's power budget.

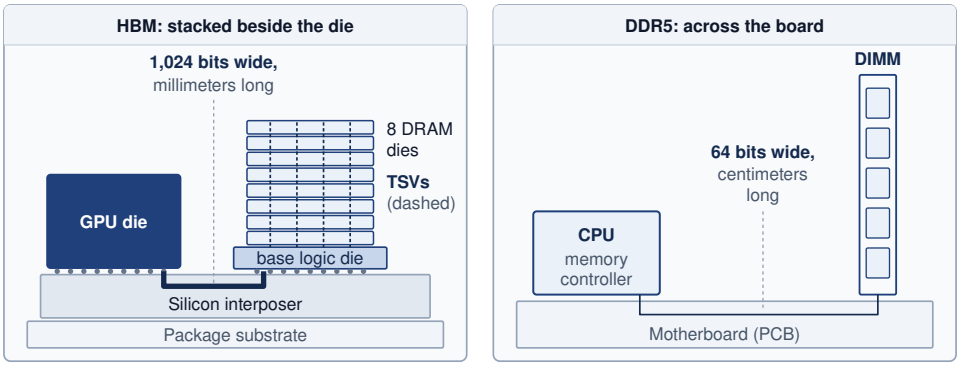

Figure 2.7: HBM 3D-Stacked Architecture: Left: high-bandwidth memory stacks DRAM dies on a base logic die, links them vertically with through-silicon vias (TSVs), and places the stack on a silicon interposer beside the GPU die, so each stack reaches the GPU over a 1,024-bit interface only millimeters long. Right: DDR5 instead sits on a DIMM across the motherboard, reached over a 64-bit channel through centimeters of board trace. The wider, shorter path raises bandwidth and lowers energy per bit, but HBM bandwidth and capacity remain finite, so the memory wall moves rather than disappears.

the short interconnects present less capacitance and need less driving current. The same reduced loading lets the 1024-bit interface run at moderate per-pin rates instead of the high-power signaling that long board traces require. Propagation delay shrinks by only a fraction of a nanosecond, small next to the DRAM array and controller delays that dominate access latency, so HBM's advantage lies in bandwidth and energy per bit rather than in latency. Modern accelerators therefore widen the memory interface across an interposer rather than driving off-package pins faster; set side by side in table 2.5, the bandwidth, interface-width, and energy differences all trace to placement and width rather than to faster signaling.

Checkpoint 2.1: HBM and the memory wall

These questions check whether the physical benefits of 3D-stacked memory are clear:

- ☐ **HBM vs. DDR5**: explain why HBM achieves higher bandwidth than DDR5 despite having a lower clock frequency.
- ☐ **Off-package signaling cost**: explain why off-package DDR memory has higher energy per bit and how vertical stacking (TSVs) reduces that cost.
- ☐ **Capacity overflow**: locate where the remaining 60 GB must reside when an H100 has 80 GB of HBM and a 70B parameter model requires 140 GB for FP16 weights.
- ☐ **Bandwidth not capacity**: classify the claim as true or false: HBM's primary benefit is increasing the total memory capacity available to the GPU.

Table 2.5: HBM vs. Standard DRAM Comparison: HBM achieves its bandwidth advantage through three simultaneous innovations: 3D die stacking (more bits per package), TSV interconnects (shorter signal paths), and on-package placement (proximity to the processor). The bandwidth figures describe each memory's local requester path: CPU to host DDR5 and H100 GPU to its own HBM3. A GPU request to host DRAM must also traverse the node interconnect and is not represented by the DDR5 channel figure. The energy row uses published per-bit access estimates for off-package DRAM at the 45 nm node (Horowitz 2014) and for HBM2 (O'Connor et al. 2017) rather than measurements of the DDR5 and HBM3 parts named in the column heads.

Metric	CPU-local host DRAM (DDR5)	GPU-local accelerator HBM (H100 HBM3)	Scaling factor
Mechanism	2D PCB traces	3D die stacking	-

Metric	CPU-local host DRAM (DDR5)	GPU-local accelerator HBM (H100 HBM3)	Scaling factor
Placement	Socketed DIMMs	On-package (substrate)	Physical proximity
Bandwidth	~200 GB/s	~3.35 TB/s	~17× faster
Interface width	64-bit	1024-bit per stack	16× wider
Energy	~20 pJ/bit	~4 pJ/bit	~5× efficiency

This bandwidth advantage comes at a price. HBM costs approximately $10–$15 per GB, compared to roughly $3 per GB for DDR5 server memory. For an H100 with 80 GB of HBM3, the memory alone represents approximately $800–$1,200 of the accelerator's manufacturing cost. For a B200 with 180 GB of HBM3e, the memory cost rises to $1,800–$2,700 per accelerator, making HBM one of the most expensive components in the system.

Packaging adds manufacturing stages beyond fabricating the memory dies. The stacking process includes thinning, alignment, bonding, and vertical interconnection, so usable stack yield depends on both the dies and their assembly. Those yields compound, because a 12-high HBM3e stack is usable only if every die and every bonding step succeeds, so even small per-step losses multiply into a complete-stack yield below that of a single DRAM die. The silicon interposer must also accommodate the processor and multiple HBM stacks. With each stack occupying roughly 100 mm^2, an illustrative large package that pairs a reticle-sized processor with several stacks needs an interposer exceeding 1,000 mm^2. Testing, repair provisions, and assembly yield all affect cost; die count alone does not determine how many complete stacks remain usable. Supply is a constraint as well as cost. HBM production is concentrated among a few manufacturers whose stacking steps need specialized equipment, so memory availability can pace accelerator delivery, a lead-time risk that section 12.9.5.1 plans for.

For our 175B model, this cost-capacity trade-off explains why accelerators typically offer 80 GB to 192 GB of HBM while the host server provides 512 GB to 2 TB of DDR. The *fast* memory holds state touched every step (weight shards, activations, gradients, and usually the sharded optimizer state), while the *cheap* memory holds staging buffers, data loading queues, and colder state that capacity pressure forces out, such as offloaded optimizer shards.

The boundary between what resides in HBM and what resides in DDR is a design parameter for training frameworks, and managing it is one of the jobs of ZeRO (Zero Redundancy Optimizer) sharding and offloading strategies. Sharding divides optimizer, gradient, and parameter state across workers; offloading places colder state in host DRAM or Non-Volatile Memory Express (NVMe) when HBM capacity is the binding limit. Getting this boundary wrong in either direction is costly. Placing too much data in HBM wastes expensive capacity, while placing too much in DDR creates bandwidth stalls that idle the arithmetic units.

↪ Memory-efficient sharding of optimizer, gradient, and parameter state is developed in section 6.3.3.

Lighthouse 2.1: Archetype B (DLRM at Scale): The capacity wall

Archetype A (GPT-4/Llama-3) has so far exposed the bandwidth side of the memory wall. Archetype B (DLRM at Scale) exposes the capacity side, the boundary between fast and cheap memory just examined. An embedding table of 10 TB or more cannot fit into the 80 GB HBM of a single H100, nor can it fit into the aggregate HBM of a single 8-GPU node (640 GB nameplate, about 687 GB in decimal bytes). Capacity alone requires at least 15 such nodes before accounting for replication, hot-table headroom, and bandwidth balance. Sharding makes lookup routing and result exchange part of the request path. An AllToAll exchange sends distinct portions to different participants; its traffic depends on the batch and placement. Clearing the capacity limit therefore creates a second question about bandwidth and latency.

2.3.1.1 Energy cost of data movement

Data movement also dominates energy. At 3.97 pJ per bit, the HBM2 access energy that O'Connor et al. estimate (2017), reading one 16-bit operand costs about 64 pJ. One FP16 multiply-accumulate

costs roughly 1.5 pJ at the 45 nm node (Horowitz 2014). The raw fetch therefore costs about 42× the arithmetic it feeds, until tiling and reuse spread it across many operations. From off-package DRAM at about 20 pJ per bit, the same fetch costs roughly 5× more again. That ratio is why accelerators devote so much silicon to reuse, and why the levers that cut bytes, such as quantization and fusion, improve power efficiency as well as speed. Section 5.2 extends the per-boundary accounting to collective communication, and section 15.4.2.3 turns these per-byte costs into an energy roofline.

2.3.1.2 HBM generations and the scaling boundary

HBM has grown alongside model sizes, but far more slowly than parameter counts, which is why capacity pressure keeps pushing models across devices. Each generation increases the number of stacked dies, the signaling rate per pin, and the total capacity per stack. The projected HBM4 generation in table 2.6 marks the scaling boundary, doubling the interface width from 1024 to 2048 bits for the first time since HBM's introduction, because pin-rate increases alone no longer sustain the required bandwidth growth. The wider bus also needs more interposer area for routing, one reason accelerator packages have grown as bandwidth targets rise.

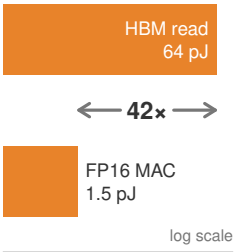

Data movement dwarfs arithmetic in energy cost.

Table 2.6: Evolution of High-Bandwidth Memory: Each generation increases aggregate accelerator memory bandwidth, tracking the growth of large model sizes. The jump to HBM4 doubles the interface width for the first time since HBM's introduction, signaling that pin-rate increases alone can no longer sustain the required bandwidth growth.

Metric	HBM2e (A100)	HBM3 (H100)	HBM3e (B200)	HBM4 (future)
Peak bandwidth	~2.04 TB/s	~3.35 TB/s	~8 TB/s	12–16 TB/s (projected)
Typical capacity	40–80 GB	80–96 GB	180 GB	288 GB+
Interface width	1024-bit	1024-bit	1024-bit	2048-bit
Stack height	8 dies	8–12 dies	12 dies	16 dies

The generations in table 2.6 move both capacity and bandwidth for our running example, and the step to HBM3e matters most. Using decimal bytes consistently, an A100 with 85.9 GB of HBM2e holds 24.5 percent of our 175B model's FP16 weights. An H100 with 85.9 GB of HBM3 holds the same fraction while delivering the data 64 percent faster. A B200 with 193.3 GB of HBM3e raises the resident share to 55.2 percent at about 8 TB/s. Because even HBM3e leaves nearly half the model outside a single package, multiple accelerators are still needed in a node (section 2.5).

Quantization changes the capacity story for inference. Storing weights in INT8 or INT4 halves or quarters the footprint, and section 2.8.1 works through how many accelerators the running model then needs. The capacity constraints that drive multi-accelerator nodes for training, where FP16 or BF16 precision is typically required, are substantially relaxed for inference, where INT8 or INT4 quantization is often acceptable, which is one reason training and inference infrastructure have different optimal configurations.

Aggregate package bandwidth scales with the number of HBM stacks that can sit beside the processor die on the silicon interposer. The H100 has 5 HBM3 stacks, each providing approximately 670 GB/s, for a total of 3.35 TB/s aggregate bandwidth. A B200-class Blackwell GPU package uses 8 HBM3e stacks for an aggregate of about 8 TB/s. The number of stacks is constrained by the interposer area available for HBM placement. Each HBM stack occupies approximately 100 mm^2 of interposer area, and the total interposer must accommodate both the processor die(s) and all HBM stacks. Larger interposers allow more HBM stacks (and therefore higher bandwidth and capacity) but are more expensive and harder to manufacture.

The interposer area constraint creates a concrete design tension. Making the processor die larger (more Tensor Cores, more SMs) leaves less interposer area for HBM stacks, potentially reducing bandwidth. Making the processor die smaller (fewer Tensor Cores) frees interposer area for more HBM but reduces peak compute throughput. The optimal balance depends on the target workload's position relative to the ridge point. Compute-bound workloads, whose arithmetic intensity sits

above it, benefit from a larger processor die (more Tensor Cores). Bandwidth-bound workloads below it benefit from more HBM stacks, so a workload's arithmetic intensity predicts which side of the balance it falls on before any silicon is committed.

2.3.1.3 Memory bandwidth and token latency

Adding or upgrading HBM stacks raises both memory *capacity* (how many gigabytes the package can store) and memory *bandwidth* (how many terabytes per second it can deliver), which answer two separate procurement questions. Capacity determines whether a model's weights *fit*, with headroom for the remaining live state; bandwidth limits how fast the model can *run* when data movement sets the critical path. Dense autoregressive decode at batch size 1 isolates the second question because each weight serves only one token before the next weight load. The next three calculations step down from the running model to Llama-3 70B. Its FP16 weights span two H100s rather than the running model's five, and the H200 comparison that follows needs a model whose serving behavior a real memory upgrade can change; the arithmetic carries over to the 175B model by scaling the bytes.

Napkin Math 2.1: The physics of token latency

Problem: A weight-only model of batch-1 decode for Llama-3 70B uses 141.2 GB of FP16 weights. To isolate the two resource demands, assume hypothetical capacity sufficient to hold them while retaining one H100-equivalent compute rate and HBM path. Which term gives the larger time estimate: arithmetic or weight transfer?

1. **Compute time**: Each token requires $2 \times 70.6 \times 10^9 = 1.41 \times 10^{11}$ floating-point operations. At 989 TFLOP/s of FP16/BF16 Tensor Core peak throughput, the arithmetic takes $T_{\text{compute}} \approx$ **0.14** ms.
2. **Memory time**: Loading 141.2 GB of weights from HBM at 3.35 TB/s takes $T_{\text{memory}} \approx$ **42.1** ms.

Systems insight: The memory term accounts for 99.7 percent of the sum of these two idealized times. That ratio identifies a bandwidth-dominated estimate; it is not a measurement of cycles spent idle. Raising compute throughput alone barely changes this estimate. A faster memory path reduces the weight-transfer floor, while actual token latency also includes the other memory traffic, communication, and execution costs omitted here.

Holding that weight load fixed and changing only the memory system shows how directly each HBM generation moves the floor. For the same 70.6B FP16 model, with 141 GB of weights, the single-path streaming floor is roughly 141 GB/2.04 TB/s = 69.2 ms at A100-class bandwidth. At H100-class bandwidth (the notebook's memory term) it is 141 GB/3.35 TB/s = 42.1 ms, and at B200-class bandwidth it is 141 GB/8 TB/s = 17.7 ms. The floor falls nearly in proportion to each bandwidth gain. As in the notebook, each case isolates one accelerator-equivalent HBM path even though the full FP16 model does not fit on one A100 or H100. These bounds screen memory bandwidth against a latency target; whether token latency improves by the same ratio depends on the other work on the request's critical path.

The same floor at trillion-parameter scale shows the limit of bandwidth gains alone, even at the projected 12–16 TB/s of an HBM4-class package. A dense 1T-parameter model requires roughly 2 TB of FP16 weight storage. Streaming those weights through one B200-class path at about 8 TB/s gives 2000 GB/8 TB/s = 250 ms for one batch-1 weight load. At the assumed 16 TB/s, that floor becomes roughly 125 ms. Neither path has been shown to hold the full model, and even the faster floor limits a single path to about eight tokens per second for one request. A responsive trillion-parameter service must therefore split its weights across many memory paths, shrink them with reduced precision, or both; which combination depends on the latency target and the quality the reduced precision preserves. The cleanest test of how far a memory upgrade moves real latency compares two accelerators that differ only in their memory system.

Napkin Math 2.2: The memory wall: H100 vs. H200

Comparing the NVIDIA H100 against the H200 isolates a memory-system upgrade for large language models (LLMs). While both chips share the same compute cores (same peak TFLOP/s), the H200 provides 1.4× higher memory bandwidth and 1.8× more HBM capacity.

This comparison uses a Llama 3 model with 70.6B parameters, sequence length 2,048, and batch size 1. Both cases retain FP16 weights and the hypothetical capacity convention of notebook 2.1: each streams the complete payload through one accelerator-equivalent HBM path, even though it cannot fit on one H100. The comparison isolates path latency rather than establishing two feasible single-device deployments.

The reported speedup comes from a serving-latency model whose inter-token latency adds two terms to the weight-streaming time, reads of the key-value cache that holds attention state for the request's earlier tokens and a fixed per-layer overhead for launching kernels, so it is not a raw bandwidth ratio alone. The weight-only estimate provides a comparison. Its memory term accounts for 99.7 percent of the summed compute and transfer times, so reducing that term suggests a gain near the HBM bandwidth ratio, 4.8 TB/s ÷ 3.35 TB/s ≈ 1.4×.

Systems insight: For the modeled configuration, the latency model predicts an H200 speedup of 1.4× over the H100 despite equal peak arithmetic throughput. Its agreement with the bandwidth ratio follows from the dominant memory term of that model, so a measured pair of systems would be the test. Changing batch size, context length, or partitioning requires evaluating the new configuration; the estimate does not establish one bottleneck for all autoregressive models.

The modeled memory-time share, 99.7 percent in this example, indicates the memory-boundedness of this workload under the stated assumptions. A workload this memory-bound gains almost nothing from more TFLOP/s and nearly proportionally from more TB/s, which is what the H200 comparison predicts, so bandwidth per dollar rather than peak arithmetic is the useful procurement metric for this regime. A measured execution timeline is still needed to determine how much transfer overlaps computation and which delays remain exposed, which keeps a hardware-screening calculation separate from a performance forecast. For training workloads, reusing weights across more tokens can instead shift the constraint toward arithmetic throughput; the Roofline Model tests the compute and bandwidth ceilings for the actual operation and batch shape before either specification becomes a procurement metric.

2.3.2 Roofline model

The token latency calculation in notebook 2.1 makes memory bandwidth the relevant procurement metric for the batch-1 workload. The Roofline Model of Williams, Waterman, and Patterson (Williams et al. 2009) tests whether that diagnosis survives a change in batch shape, precision, or accelerator by relating the workload's arithmetic intensity to the candidate accelerator's compute and bandwidth ceilings.

For an H100 (989 TFLOP/s FP16, 3.35 TB/s HBM), the ridge point is about 295 FLOP/byte. Most LLM inference operators fall well below this threshold, which is why the roofline remains the first diagnostic tool for identifying whether more compute or more bandwidth will improve performance. The quantity that places a workload on the plot is the arithmetic intensity introduced with tiling above, written I in FLOP/byte.

↪ Placing measured kernels on the roofline by arithmetic intensity is developed in section 9.6.3.

As the single-machine recap in section A.3.1 derives, the Roofline Model expresses the maximum achievable performance of a workload as the lesser of two ceilings:

$$\text{Achievable FLOP/s} = \min\left(R_{\text{peak}},\ \text{BW} \times I\right) \tag{2.1}$$

In equation 2.1, below the ridge point the sloped bandwidth ceiling, which grows with I, is the lower one; above it the flat compute roof binds, and a measured kernel can sit below either.

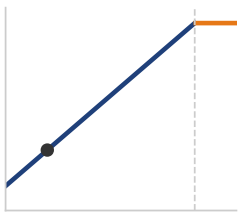

Decode sits deep in the memory-bound regime, far below the ridge.

Definition 2.3: Ridge point

Ridge point is the ML accelerator arithmetic intensity where the memory bandwidth ceiling meets the compute ceiling ($I_{\text{ridge}} = \frac{R_{\text{peak}}}{\text{BW}}$).

1. **Significance**: It is the accelerator's efficiency threshold. Below the ridge point, bandwidth supplies the lower modeled ceiling (BW); above it, arithmetic supplies the lower ceiling (R_{peak}).
2. **Distinction**: Unlike peak FLOP/s, which describes only the horizontal ceiling, the ridge point describes the balance between arithmetic and memory delivery in the architecture. A rising ridge point over hardware generations indicates that compute is growing faster than bandwidth, making utilization harder.
3. **Common pitfall**: A frequent misconception is that an accelerator has a single ridge point. In reality, it varies by precision. Because R_{peak} is higher for INT8 than FP32 while BW is constant, the ridge point for INT8 is much higher, requiring more data reuse to saturate the hardware.

Machine learning workloads land in different regions of the roofline plot depending on the operation and the batch configuration, not on the model family. Single-request autoregressive decode streams the full weight tensor (two bytes per parameter in FP16) from HBM to perform two floating-point operations per parameter, an arithmetic intensity of about 1 FLOP/byte, far below the H100's ridge point of 295 FLOP/byte. The token-latency calculation in notebook 2.1 showed that the arithmetic finishes in microseconds while the weight movement takes milliseconds. Every other regime in table 2.7, including prefill, the phase that processes a request's prompt before decode begins, comes from reusing each byte fetched from HBM across more operations. Decode also reads the key-value (KV) cache, the stored attention keys and values of earlier tokens, and because each request has its own cache, batching does not amortize that traffic.

Table 2.7: Roofline Placement by Workload Phase: Language-model intensities count only weight traffic at the HBM level, the same idealization used for figure 2.8, against an H100 FP16 ridge point of 295 FLOP/byte; the convolution value is a representative scenario figure rather than a derived count. Attention is not plotted in the figure because its placement depends on shape. Measured placement must also count activation, key-value cache, and optimizer traffic.

Workload phase	Source of reuse	Idealized intensity	What moves it
Decode, batch 1	None; each weight serves one token	1 FLOP/byte	Batching requests, lower-precision weights
Decode, batch 32	Each weight read serves every request in the batch	32 FLOP/byte	Batch size, until key-value cache traffic dominates
Prefill, 512-token prompt	Each weight read serves every prompt token	512 FLOP/byte	Prompt length and batch size
Convolution training (ResNet-50)	Filters reused across spatial positions and the batch	200 FLOP/byte	Channel depth, tensor layout, operator fusion
LLM training, 2048 local tokens	Each weight read serves every token on the device	2,048 FLOP/byte	Local microbatch size and sequence length
Attention	Keys and values reused across queries	Shape-dependent, on either side of the ridge	Sequence length, head dimension, SRAM tiling

The table's intensities are idealized reuse counts that classify an execution phase and batch shape, not a model family, and tensor layout, fusion, precision, and cache residency move the measured point. Which kernel-level change then helps, more reuse and fewer bytes below the ridge or a cheaper operation above it, is a diagnosis that section 9.6.3 develops.

A measured workload point below its roofline ceiling needs further diagnosis. Dependencies, insufficient parallel work, irregular accesses, synchronization, and power limits can all prevent a

kernel from approaching that ceiling. The gap identifies a question for profiling; it does not establish that software optimization can recover the entire difference.

Whole phases spread even wider than individual operators. The idealized intensities of decode and training differ by 2,048×, which is why the same hardware shows widely different utilization across them. Batch-1 decode sits near 1 FLOP/byte, far left of the ridge and deeply bound by memory bandwidth, while large-batch training near 2,048 FLOP/byte crosses the ridge and reaches the compute roof (figure 2.8). The training point assumes that 2048 tokens on one accelerator share each weight read, the local reuse that a large global batch alone does not guarantee, so it counts tokens rather than sequences and local rather than global batch. It is a contrast case, not a universal intensity for training. Actual placement requires counting bytes at the memory level being modeled and accounting for tensor shapes, fusion, and reuse.

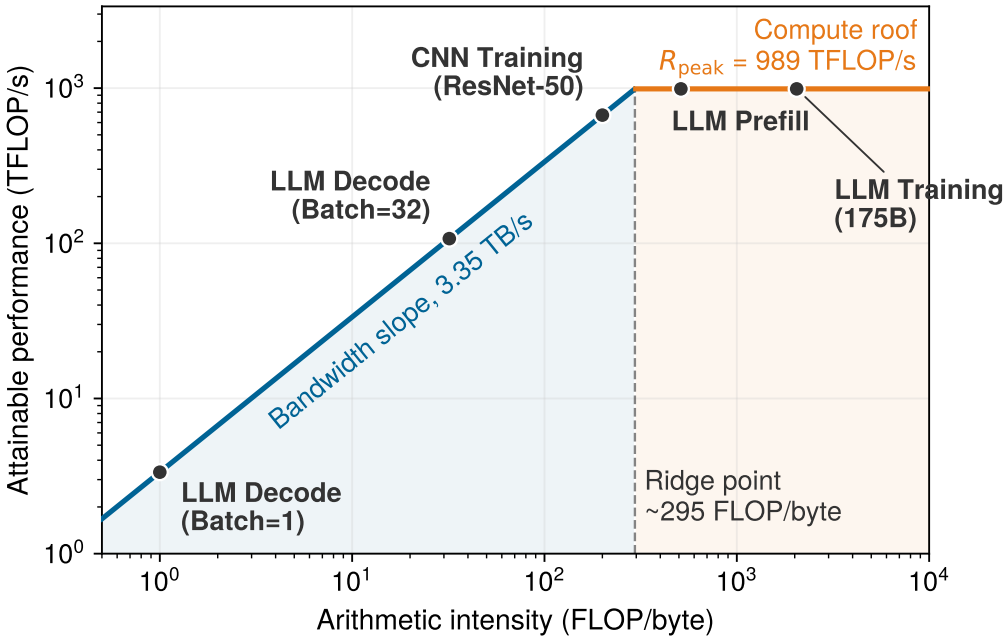

Figure 2.8: The H100 Roofline Landscape: Attainable FP16 throughput is the lower of two ceilings, the memory-bandwidth slope (blue) and the compute roof (orange), which meet at the dashed ridge point. Each dot places a representative workload by its arithmetic intensity; workloads left of the ridge stay capped by bandwidth however much compute the chip offers, and only those to its right can approach peak throughput.

The table's batching rows can be turned into a utilization bound. For inference serving, batching multiple requests together moves the workload point rightward on the plot, because each weight read is multiplied against more activations. Under the weight-traffic-only bound, batch-1 decode on an H100 can use at most 0.34 percent of the FP16 peak. At batch size 64, each weight read from HBM serves every request in the batch, the intensity rises to 64 FLOP/byte, and the bound climbs to about 22 percent of peak, a 64× gain that still leaves the workload below the ridge. The linear growth holds only in this approximation; key-value cache traffic, which batching does not amortize, activation traffic, and capacity limits lower both figures in practice. The gain has a cost. Waiting to assemble a batch can introduce queuing latency, although a serving policy can instead dispatch a smaller batch when a latency deadline approaches. Serving systems must carefully balance this trade-off between throughput (larger batches, higher utilization) and latency (smaller batches, faster response per request).

For fleet procurement, the roofline is a screening test rather than a throughput forecast. It rules out upgrades that raise a nonbinding ceiling, while the gap below the relevant ceiling identifies work still needed in software or communication. For the running model, the two ends of the landscape give opposite hardware answers. Memory capacity and bandwidth screen out infeasible accelerator

choices, while for the candidates that remain, fleet sizing also needs a sustained compute rate at a precision the workload can use, within the node's memory budget, the interconnect, and the facility power envelope.

Example 2.2: Roofline analysis: Training vs. inference

Scenario: An engineering team evaluates hardware choices for serving (decode batch size 1) vs. training (batch size 2048) a 175B parameter model on an H100 GPU (ridge point 295.2 FLOP/byte).

Diagnosis: The weight-traffic model gives single-token decode arithmetic intensity $I = 1$ FLOP/byte $\ll$ 295.2 FLOP/byte, bounding throughput to 0.34 percent of the selected peak FLOP/s. The assumed training reuse raises intensity to $I = 2{,}048$ FLOP/byte $\gg$ 295.2 FLOP/byte, making arithmetic the lower roofline ceiling. This placement identifies the relevant bound without establishing achieved throughput.

Systems lesson: Hardware selection starts with the phase's arithmetic intensity and memory footprint. The modeled training phase benefits from arithmetic capacity, while batch-1 decode benefits from bandwidth. Sustained throughput at the required quality and latency determines whether either advantage justifies its cost.

Checkpoint 2.2: Roofline diagnosis

A team is serving a 13B-parameter model at batch size 1 on an H100 and observing 25 ms per token. The profile shows low tensor-core utilization during decode.

- ☐ **Arithmetic intensity**: estimate the arithmetic intensity of the workload.
- ☐ **Peak-throughput upgrade**: explain what the roofline estimate can establish about a higher peak-TFLOP/s accelerator and what requires measured execution.
- ☐ **Binding constraint**: identify the profile measurements needed to distinguish exposed HBM transfers from communication, launch overhead, or insufficient parallel work before selecting an improvement.

2.4 Sustained Compute and Precision

A fleet sized from advertised matrix throughput can miss its training deadline when operations cannot use the advertised precision or keep the matrix units of section 2.2.5 busy. Sizing a fleet therefore requires knowing both which matrix operations are eligible for the advertised rate and how much throughput those operations sustain across the complete workload.

Tensor Core precision support widened with each generation traced in section 2.2.3, from FP16 on Volta to FP8 on Hopper, because neural network training and inference tolerate lower numerical precision than was once assumed. Transformer Engine recipes and FP8 autocast decide which supported operations run in FP8, recording tensor ranges and applying scale factors so values fit the format, but they do not establish that a layer is numerically safe at that precision. Precision policy therefore remains a software configuration and validation decision, which section 6.4.8 develops for training.

The relationship between numerical precision and hardware throughput matters for infrastructure planning because it directly affects how much useful work each accelerator can perform per second. Lower precision representations use fewer bits per number, which has two compounding benefits: more numbers fit in the same HBM capacity (reducing memory pressure), and the arithmetic units can process more operations per cycle (increasing throughput). Within a format's bits, the exponent sets the range of values it can represent and the mantissa sets how finely it resolves them. Each format an H100-class accelerator supports unlocks a different dense peak rate and plays a usual role in training or inference (table 2.8).

Table 2.8: Precision Formats and Hardware Throughput on H100: The table uses dense peak throughput where Tensor Core support exists; structured sparsity can roughly double several of these advertised peaks. The choice of precision affects both the training dynamics (convergence, numerical stability) and the infrastructure requirements (memory capacity, bandwidth utilization). TF32 preserves the FP32 exponent width but reduces mantissa precision so Tensor Cores can accelerate common FP32 training workloads.

Format	Bits	Exponent	Mantissa	H100 dense peak	Typical use
FP32	32	8	23	~67 TFLOP/s	Master weights, loss computation
TF32	19	8	10	494.5 TFLOP/s	Training (NVIDIA default)
BF16	16	8	7	989 TFLOP/s	Training, fine-tuning
FP16	16	5	10	989 TFLOP/s	Training with loss scaling
FP8 (E4M3)	8	4	3	1,979 TFLOP/s	Forward pass, weights
FP8 (E5M2)	8	5	2	1,979 TFLOP/s	Backward pass, gradients
INT8	8	N/A	N/A	1,979 TOPS	Post-training quantization

For training, many workloads use the standard mixed-precision recipe, which decouples storage from arithmetic precision by keeping FP32 master weights while the forward and backward passes run in BF16 or FP16 at full Tensor Core throughput. FP16 training relies on loss scaling, which multiplies the loss so that small gradients stay within FP16's narrow range. The node memory examples in section 2.5.3 assume retained FP16 gradients; implementations that retain a full FP32 accumulation buffer must add that buffer to their peak budget. The FP8 recipes described earlier extend the same decoupling one step further, running eligible matrix operations in eight bits while the master weights stay in FP32.

The practical implication for fleet design is that peak TFLOP/s specifications are precision-dependent. The H100 delivers 989 TFLOP/s in FP16 but twice that (1,979 TFLOP/s) in FP8. FP8 therefore doubles the advertised arithmetic ceiling for supported operations on the same hardware. The complete training run gains less when communication, memory traffic, or higher-precision operations remain on the critical path. Precision engineering is therefore a primary optimization lever for infrastructure planners, alongside its effect on model accuracy.

↪ The byte widths of each format and the per-parameter memory of mixed-precision training are collected for reference in section A.6.

Approaching peak throughput requires pipeline balance. Operand staging and reuse must supply the arithmetic units at the rate the operation requires. Kernel fusion and tiled attention such as FlashAttention improve that balance by keeping intermediate values on chip instead of making separate HBM round-trips, mechanisms that section A.5 collects for reference. For our 175B model, the fleet-level gain depends on how much step time the affected kernels consume. A faster attention kernel shortens the step only in proportion to attention's share of it, the same Amdahl accounting that bounds every partial speedup.

2.4.1 Peak vs. sustained throughput

A training run budgeted for 1,000 GPU-hours from datasheet TFLOP/s will consume 2,500 GPU-hours in practice when the workload sustains only 40 percent of peak throughput. Infrastructure planning therefore separates *peak* throughput (the maximum rate for the vendor-specified operation and precision) from *sustained* throughput (what the complete workload delivers). Peak TFLOP/s assumes that eligible arithmetic units remain fed under the stated sparsity, precision, and clock conditions; it does not include memory stalls, ineligible operations, communication, or software overhead. Sustained fractions therefore vary with model shape, batch size, precision, implementation, and scale. The gap arises from several sources, each representing a different physical or software limitation.

Memory stalls occur when the Tensor Cores complete their current tile multiplication before the next tile has been loaded from HBM. Even with the H100's 3.35 TB/s HBM bandwidth, the Tensor

Cores can consume data faster than HBM can deliver it for certain matrix dimensions, creating idle cycles while the arithmetic units wait for data.

Dependency stalls arise when one operation must complete before the next can begin, leaving some processing elements idle. In a Transformer layer, the attention computation must complete before the feed-forward network can begin, and the feed-forward computation must complete before the next layer's attention can start. These sequential dependencies create brief periods where some hardware resources are unused, even within one accelerator. When the model is split into sequential stages on different accelerators (pipeline parallelism), the same dependencies become pipeline bubbles, idle slots in which a whole stage waits for its predecessor. Splitting each batch into microbatches reduces the waste by feeding later microbatches into stages that would otherwise wait. The term pipeline bubble refers only to this cross-device case, whose scheduling chapter 6 develops. At the infrastructure level, the lesson is that peak throughput assumes every unit stays busy, while real training includes dependency gaps that lower sustained utilization.

AllReduce operations used in tensor parallelism and gradient synchronization used in data parallelism can likewise delay computations that depend on their results; even when communication overlaps computation on separate streams, operations that depend on the collective output cannot proceed until the transfer completes. Kernel submission, memory allocation, and framework control also consume time, and when each kernel completes in microseconds, launch overhead of a similar size becomes a comparable share of the step that only a host-and-device execution trace can separate from slow arithmetic or memory access. Achieving high sustained utilization is therefore largely a matter of minimizing the time Tensor Cores spend idle between operations, which is why kernel fusion and specialized kernel libraries such as FlashAttention and the Transformer Engine exist.

The last source of the gap is the device itself, and no kernel change removes it. Device variation, sometimes called the silicon lottery, means that nominally identical accelerators sustain different clocks. Clock behavior also depends on the configured power limit, temperature, and workload, so a slow participant should not be diagnosed from its product label or frequency alone. In synchronous training, a participant that finishes dependent work late delays its whole process group, the set of processes that run a collective together. One device sustaining a clock about 2.5 percent below its peers can slow every synchronous step it joins by up to that fraction, because the other participants wait for it at each collective. Fleet management systems therefore track per-device sustained throughput and can move persistent underperformers to pools, such as inference serving, where they do not gate a synchronized group. Profiling must determine whether local arithmetic, memory, communication, or a power limit caused that delay before changing placement or replacing hardware.

Capacity planning must use a measured or conservatively modeled sustained rate, not the peak specification. A run estimated at 1,000 GPU-hours from peak FLOP/s requires $1/\eta_{\text{hw}}$ times that estimate when its measured utilization fraction is η_{hw}. Infrastructure teams track model FLOPs utilization (MFU), the useful model computation per second divided by the selected hardware peak, as the standard measure of η_{hw}, but comparisons remain meaningful only when model-FLOP conventions, precision, hardware peak, and system scale match.

↪ MFU as the measure of program goodput, one factor of fleet-level ML goodput, is developed in section 9.2.2.

The accelerator's physics establish a concrete problem. Our 175B model requires 350 GB of memory for its weights alone in FP16, and Adam-style training state raises the requirement to multiple terabytes once optimizer moments, gradients, and activations are included. A single H100 provides 80 GB of HBM. No single accelerator can hold this model, an instance of the memory capacity gap (principle 2), in which frontier memory requirements outgrow single-accelerator HBM. The system must expand to the next physical level: the node.

2.5 The Node

The running model needs more memory than one accelerator provides. A node supplies several accelerators in one chassis, but each retains its own memory; dividing the work creates communication that the local fabric must sustain. Tensor parallelism divides an operation across accelerators, pipeline parallelism assigns successive model stages to different groups, and data parallelism assigns different training examples to model replicas whose updates must be combined. These strategies impose different traffic patterns, and the node's internal fabric decides which of them it can carry at the speed the partition needs.

Aggregating accelerators solves the capacity problem only if the interconnect between them is fast enough that the group behaves as one machine rather than as several machines passing messages. The 175B model's weights span at least five accelerators, and the full training state spans an entire 8-H100 node and more, so some model state must live on different chips. Depending on the partition, execution exchanges partial results, activations, gradients, or parameter shards. Whether those exchanges are cheap or ruinous depends on the local fabric, which is why the node is the first scale level where the accelerator's specifications matter less than the wires between accelerators. The training system can combine tensor parallelism, optimizer sharding, and memory offload as needed to meet its capacity budget, and each strategy places a different demand on the interconnect.

Definition 2.4: Node

Node is the ML training cluster's physical server chassis that aggregates multiple accelerators (typically 8) through a high-speed intra-node interconnect (NVLink or ICI), providing a physical packaging and failure boundary. The fast communication domain can coincide with one chassis or extend across several, depending on the fabric.

1. **Significance**: The NVLink bandwidth within a DGX H100 node reaches 900 GB/s bidirectional, or 450 GB/s per direction, while inter-node InfiniBand NDR provides 50 GB/s, a 9× per-direction gap that constrains where each parallelism strategy can run. For this DGX H100 topology, frequent tensor-parallel collectives favor the intra-node NVLink domain; pipeline parallelism can span nodes when the network sustains stage-boundary activation traffic. The node is also the primary failure domain. A single node failure in a bulk synchronous parallel training job, where workers advance in lockstep at step barriers, idles the entire cluster until recovery.
2. **Distinction**: Unlike a single accelerator (which provides fast HBM bandwidth but limited capacity), a node aggregates 8× the HBM capacity and 8× the compute of a single chip. That is enough to place large FP16 weight shards (8×80 GB = 640 GB on DGX H100), but not enough to keep the full Adam training state in HBM.
3. **Common pitfall**: A frequent misconception is that nodes in a distributed training job operate independently. In bulk synchronous parallel training, all nodes are synchronously coupled at every step barrier. A performance degradation on one node (from thermal throttling, a faulty NIC, or a slow storage read) stalls the entire cluster and drops overall MFU for the duration of that step.

Holding the overflow in host DRAM cannot replace wiring eight accelerators together. A server's DDR5 channels supply ample capacity (512 GB to 2 TB per node) at a fraction of HBM's cost, but their bandwidth is orders of magnitude too slow to feed Tensor Cores on the active forward and backward path. Pooling the HBM of eight accelerators inside one chassis keeps the active working set in high-bandwidth memory, with the caveat that each accelerator's HBM remains private to its own package. Software must shard or replicate tensors explicitly and invoke collectives or peer-to-peer transfers whenever an operation crosses device boundaries, so an eight-GPU node delivers 8× aggregate capacity and 8× aggregate compute without ever behaving like a single shared-memory processor.

The capacity arithmetic also explains the appeal of multiple packages. Holding a 175B model's 350 GB of FP16 weights on one accelerator, at the example 17–24 GB per HBM stack, would require 350 GB spread across roughly 15–21 stacks. That exceeds the 5-stack H100 and 8-stack B200-class configurations used here. Such a package would need its own area, yield, power, and cooling analysis. Distributing capacity across 8 accelerators instead uses available packages and pays for communication between their private memories.

2.5.1 Bandwidth hierarchy

Placing a layer-by-layer tensor-parallel collective across a 50 GB/s inter-node link instead of a 450 GB/s intra-node link multiplies its transfer floor ninefold and stalls every accelerator in the group. Multi-accelerator systems are governed less by the *number of chips* than by the *speed at which they can*

exchange data across physical boundaries. Data movement speed drops by orders of magnitude as it crosses package, board, and rack boundaries, each with a different physical medium, connector technology, and signaling constraint. This physical ordering of data transfer rates, from on-chip SRAM at the fast end to the wide-area network at the slow end, is the bandwidth hierarchy; this ordering dictates where in the hierarchy each type of parallelism can operate efficiently. The cliff between intra-node and inter-node links is where the communication wall first appears.

The hierarchy reflects the physical cost of distance. Shorter signal paths over denser wiring achieve higher throughput at lower energy per bit, so at each boundary the effective bandwidth (BW) drops by roughly an order of magnitude while latency (L_{lat}) rises. The exact size of each gap depends on link count, signaling, and the measurement boundary, but the gradient itself sets a scaling ceiling for distributed training. The common mistake is to treat all cluster communication as equal. Parallelism placement must be hierarchy-aware, because putting high-frequency synchronization such as tensor parallelism on a slow tier leaves the arithmetic units (R_{peak}) waiting and collapses scaling efficiency (η_{scaling}), the share of ideal linear speedup that the cluster delivers.

The bandwidth hierarchy can be read as concentric zones around each accelerator (figure 2.9), losing 7.4× to 9× of bandwidth at each physical boundary on the way out through NVLink, InfiniBand, and remote storage. The innermost zone is the HBM on the same package, with terabytes per second of bandwidth and sub-microsecond access latency. The next zone holds the other accelerators in the node's NVLink domain, reachable at hundreds of gigabytes per second per GPU with microsecond-scale transfer latency. The outermost zone spans the other nodes in the cluster, reached through InfiniBand ports that each supply tens of gigabytes per second per direction, with point-to-point latency measured in single-digit microseconds. Latency figures are comparable only when they measure the same thing (a memory access, a link transfer, or a complete collective), and bandwidth ratios only when they compare matching endpoints and directions, as the representative paths of table 2.9 do.

Table 2.9: The Bandwidth Hierarchy: These representative paths have different bandwidth budgets. Link bandwidths are per direction; link count, signaling, and system configuration determine the gap between tiers. These example bandwidth gaps favor placing frequent tensor-parallel collectives within the fast communication domain, which need not coincide with a server boundary.

Domain	Interconnect	Bandwidth	Latency	Scaling limit
Intra-package	Silicon interposer	~3.35 TB/s	$<$ 100 ns	Single chip
Intra-node	NVLink/ICI	~450 GB/s	~0.5 µs	Node (8–16 chips)
Intra-node (I/O)	PCIe Gen5 x16	~64 GB/s	~1 µs	CPU-GPU, NIC-GPU
Inter-node	InfiniBand NDR	~50 GB/s	~5 µs	Pod (thousands)

Napkin Math 2.3: The physics of the staircase

Problem: A 10 GB buffer must move one hop at each tier of the hierarchy, with a local HBM read as the reference point. How much does the tier it crosses change the transfer time?

Math: Transfer time is $T_{\text{transfer}} = D_{\text{vol}}/\mathrm{BW}$.

1. **HBM (local read, the reference point)**: 10 GB / 3,350 GB/s ≈ 3.0 ms.
2. **NVLink (intra-node)**: 10 GB / 450 GB/s ≈ 22.2 ms.
3. **InfiniBand (inter-node)**: 10 GB / 50 GB/s ≈ 200 ms.

Systems insight: For the same payload, the example inter-node path has a transfer-time lower bound 9× larger than the example intra-node path. Size alone does not decide which traffic may cross that cliff; frequency and overlap do. Tensor parallelism exchanges activations inside every layer, on the critical path, so it favors the fast local NVLink domain in this configuration. Data

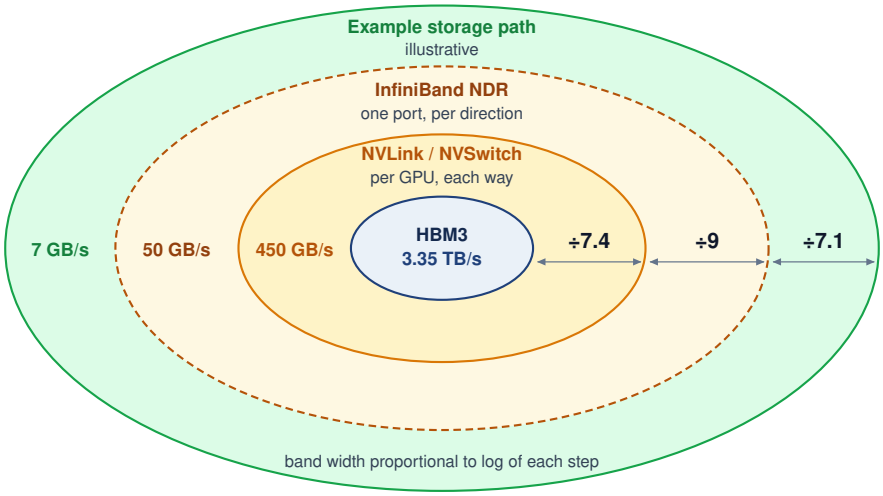

Figure 2.9: The Infrastructure Bandwidth Hierarchy: Concentric zones place each tier an ML workload crosses farther from the accelerator's own memory. HBM (3.35 TB/s) sits at the center; the outward tiers are NVLink/NVSwitch (450 GB/s per direction per GPU), one InfiniBand NDR port (50 GB/s per direction), and an illustrative 7 GB/s storage path. Each step outward cuts bandwidth by roughly 7.4, 9, and 7.1 times, not by a universal tenfold law. Exploiting this hierarchy through tiered algorithms is the primary task of distributed systems engineering.

> parallelism moves the largest buffer, the full gradient, but only once per step and overlapped with the backward pass, so it can cross the scarce InfiniBand tier. A partition that fits in memory can still spend much of a step waiting if its exposed exchanges exceed the communication budget.

Each cliff in the bandwidth staircase has a physical cause at its distance scale. Millimeter-long traces on a silicon interposer can carry a wide, slow-clocked parallel interface without serializer/deserializer (SerDes) transceivers. NVLink crosses tens of centimeters of board copper, which requires SerDes transceivers with equalization to recover the signal at high per-lane rates. PCIe pays for universal compatibility with protocol overhead (packet headers, flow-control credits, and error checking). InfiniBand crosses meters of cable between racks, adding optical conversion on longer runs, forward error correction, and switch hops. Section 3.3.3 develops the signaling physics behind these cliffs. For node design, link count and sustained bandwidth must be evaluated together with the power and cooling load of the components that provide them.

The practical consequence for model partitioning is that tensor parallelism (TP), which splits individual matrix multiplications across accelerators and requires collective exchanges within layers, generates communication volume proportional to the model's hidden dimension hundreds of times per second. Over NVLink, each of these exchanges takes under a millisecond for the running model (section 2.5.2 computes the idealized lower bound). At the reference InfiniBand port rate, the bandwidth-only lower bound for the same exchange is 9× as large; the step-time effect also depends on startup, routing, and overlap. For this topology TP therefore stays within a node, although a system whose fast scale-up fabric spans several chassis can place a tensor-parallel group across them; communication cost, rather than the server enclosure, determines the useful partition.

The hierarchy also explains why pipeline parallelism occupies an intermediate position in the bandwidth requirements. In pipeline parallelism, the model is divided into sequential stages, with each stage assigned to a different group of accelerators. The communication between stages consists of activations flowing forward during the forward pass and gradients flowing backward during the backward pass, with a volume proportional to the batch size times the hidden dimension.

For our 175B model with hidden dimension 12,288 and one sequence of 2,048 tokens per microbatch, the activation tensor at each stage boundary is approximately $1 \times 2,048 \times 12,288 \times 2B$ bytes (FP16) $\approx$ 50 MB. This is far smaller than the 350 GB gradient AllReduce required by data parallelism, which is why pipeline parallelism places less demanding requirements on the inter-node network.

Activation transfers occur once per microbatch per stage boundary. That is far less often than tensor parallelism's AllReduce inside every layer, yet more often than data parallelism's gradient exchange, which happens once per step. Small messages at an intermediate frequency let pipeline parallelism tolerate the lower bandwidth of inter-node links, which makes it the natural way to split a model whose layers do not fit in one node.

Both pipeline and data parallelism therefore cross node boundaries, for different reasons (table 2.10). Pipeline traffic is small per message and sits between the other two in frequency. The data-parallel gradient exchange is the largest transfer, but it occurs once per step and can overlap the backward pass. That overlap makes it the most latency-tolerant of the three, although any unfinished exchange still delays the update. For the example topology, the resulting mapping places tensor parallelism within the node, pipeline stages across nodes, and data-parallel groups across the cluster. This *de facto* pattern is a starting point for the capacity and timing calculation; a different scale-up domain, message size, or overlap opportunity can justify a different placement.

Table 2.10: Traffic Shape by Parallelism Strategy: The placement rule follows from size and frequency together. Frequent exchanges on the critical path need the fastest tier, while the largest exchange can cross the slowest tier because it happens once per step and overlaps the backward pass.

Strategy	Size per exchange	Frequency	Overlap with computation	Placement on this topology
Tensor parallelism	Activation-sized partial results	Inside every layer, forward and backward	Little; later operations depend on the result	Within the NVLink domain
Pipeline parallelism	One microbatch's activations at a stage boundary	Once per microbatch per boundary	Partial, through microbatch scheduling	Across nodes, adjacent stages close together
Data parallelism	The full gradient	Once per step	Substantial, during the backward pass	Across the cluster

The mapping constrains how the scheduler assigns devices and nodes. Members of a tensor-parallel group belong inside one fast domain (on this topology, one chassis joined by NVLink), because their dependent exchanges recur inside every layer. Adjacent pipeline stages belong in the same rack or neighboring racks, where their activation and gradient traffic crosses the fewest InfiniBand switch hops. Data-parallel groups can span the whole cluster. Physical adjacency helps only when it corresponds to better connectivity. The node's PCIe hierarchy adds another placement constraint by connecting accelerators to host CPUs, network interface cards (NICs), and storage, including the path used to reach host DRAM.

2.5.1.1 The PCIe hierarchy within a node

While NVLink carries GPU-to-GPU traffic, PCIe Gen5 connects accelerators, host CPUs, network adapters, and storage through the node's peripheral hierarchy. Host DRAM attaches to the CPU's memory controllers; a GPU transfer to that memory also traverses its host-facing PCIe path. PCIe bandwidth, approximately 64 GB/s per direction (128 GB/s bidirectional) per x16 link, can therefore constrain data loading, checkpointing, or GPU-to-network transfers. Which traffic competes depends on the actual links and switches those paths share.

A standard DGX H100 node features two host CPUs, each managing a PCIe root complex, the CPU-side hub where its PCIe lanes originate, and PCIe switches fan those lanes out. Each of the 8 GPUs and each of the 8 InfiniBand host channel adapters (HCAs, the network interface cards of an InfiniBand fabric) has its own x16 link (figure 2.10). This configuration provides an aggregate theoretical bandwidth of about 2 TB/s within the chassis. However, this bandwidth is shared among competing traffic streams. During a single training step, the PCIe bus simultaneously carries host-to-GPU data batches, GPU-to-HCA gradient shards for inter-node AllReduce, and periodic CPU-to-NVMe checkpoint writes. The gradient shards travel through GPUDirect remote direct

memory access (RDMA), which moves data between GPU memory and the HCA without a host-memory copy. Contention between these streams can cause unexpected pipeline stalls, particularly when gradient synchronization saturates the links typically used for data loading.

The dual-CPU architecture introduces significant non-uniform memory access (NUMA) effects. The 8 GPUs are physically partitioned, with 4 connected to CPU 0's root complex and 4 to CPU 1's. A GPU connected to CPU 0 can reach CPU 0's DRAM through its local PCIe root complex, whereas access to CPU 1's DRAM also traverses the inter-socket link. The latency and bandwidth penalty depends on the processor, socket interconnect, PCIe topology, transfer size, and concurrent traffic, so no fixed percentage applies across nodes. Data-loading pipelines should measure this topology and place worker processes, pinned buffers (host memory locked in place so that devices can read it directly), NICs, and target GPUs within the same NUMA domain when the local path improves sustained ingestion. Placement is therefore part of the parallelism strategy, because a scheduler that ignores PCIe roots, NIC locality, and NUMA domains can turn an otherwise valid tensor- or pipeline-parallel plan into a host-bound bottleneck.

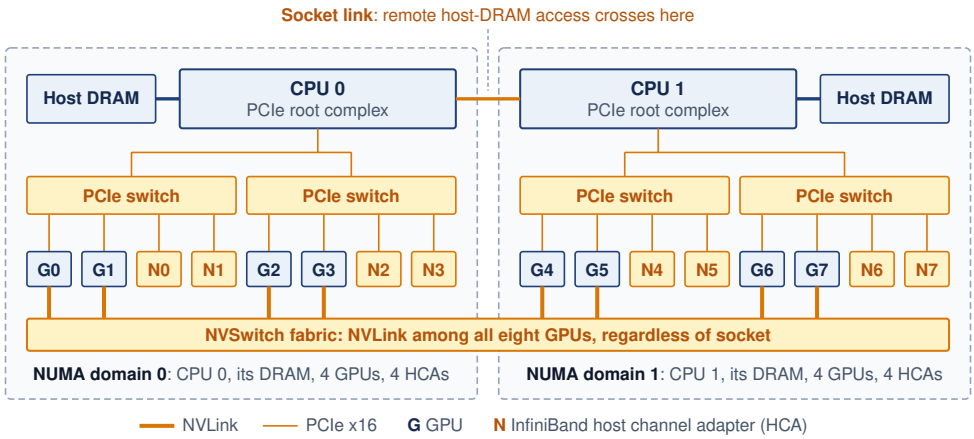

Figure 2.10: Two Fabrics inside One Eight-GPU Node: NVSwitch joins all eight GPUs for GPU-to-GPU traffic, but every path to host memory, storage, or the network runs through a PCIe tree rooted at one of two CPU sockets. A GPU, its InfiniBand host channel adapter (HCA), and the host DRAM it reaches most cheaply therefore sit in one non-uniform memory access (NUMA) domain, and a transfer to the other socket's DRAM also crosses the socket link. The two PCIe switches per socket are schematic; switch count and lane routing vary by server design.

2.5.2 Dense node designs

↪ Mapping data, tensor, and pipeline parallelism onto intra-node and inter-node bandwidth tiers is developed in section 6.6.

Given the bandwidth hierarchy, the node must connect 8 accelerators with enough bandwidth for tensor-parallel operations across their private memories. A switched fabric can provide full bisection bandwidth, enough capacity across any equal division of the endpoints to carry their full injection rate. Comparing a ring, a mesh, and a switched design shows what connectivity costs and which communication patterns it supports.

Eight GPUs in one chassis can be wired as a ring, a full mesh, or a switched fabric (figure 2.11). A ring is inexpensive, needing only G_{node} links for G_{node} GPUs, and suits ring-AllReduce, in which each GPU sends to one neighbor while receiving from the other so that every link stays busy. Tensor parallelism, however, issues frequent collectives after sharded matrix operations. These include ReduceScatter, which leaves each participant with one shard of a combined result, and AllGather, which assembles shards at every participant. A ring's fixed multi-hop paths (up to four hops between distant GPUs) constrain those patterns. A full mesh reaches every peer in one hop but needs $G_{\text{node}}(G_{\text{node}}-1)/2 = 28$ links for 8 GPUs, a count that grows quadratically and soon becomes impractical to route. The NVSwitch crossbar puts every pair one switch hop apart and provides full-bisection switched connectivity with a manageable number of switch chips.

Consider the NVIDIA DGX H100 architecture. Eight H100 GPUs sit on a single baseboard, each connected to four NVSwitch chips via 18 NVLink 4.0 lanes. The NVSwitch chips form a non-blocking

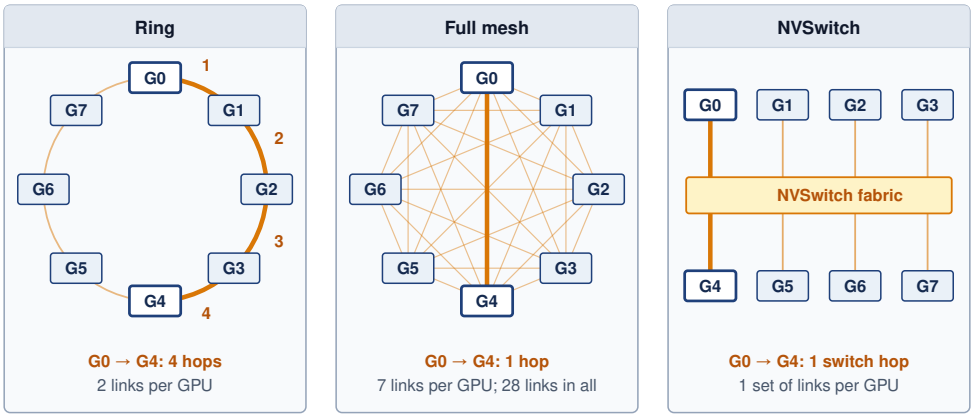

Figure 2.11: Intra-node GPU Topology Comparison: Three ways to wire eight GPUs inside one node. Left: a ring gives each GPU only two links, so traffic between distant GPUs passes through up to four hops. Center: a full mesh reaches every peer in one hop but needs $G_{\text{node}} - 1$ links per GPU and $G_{\text{node}}(G_{\text{node}} - 1)/2$ links in total, which becomes impractical to route beyond about eight GPUs. Right: an NVSwitch fabric gives each GPU one set of links to the switch and puts every pair one switch hop apart, with up to 450 GB/s per direction per H100 SXM GPU (900 GB/s bidirectional); that per-GPU injection bandwidth is not the aggregate bisection bandwidth of the switch fabric.

crossbar[5] through which any GPU can reach any other GPU. Each GPU's port budget is 900 GB/s bidirectional, or 450 GB/s in one direction, shared across that GPU's simultaneous transfers; it is not a separate full-bandwidth link for every GPU pair.

The result is a node with 640 GB of aggregate HBM capacity and a combined peak compute throughput of 8 × 1,979 TFLOP/s, while each GPU retains private HBM. Training frameworks partition tensors explicitly across those devices, and NVSwitch reduces the communication cost between shards without making placement or collective synchronization transparent. The NVSwitch chips themselves consume approximately 200 W each (800 W total for the NVSwitch fabric), a nontrivial power overhead justified by the bisection bandwidth it provides. In physical terms, the NVSwitch fabric alone consumes more power than an entire previous-generation GPU (the V100 drew 300 W).

Tracing one microbatch of the running model through its tensor-parallel exchanges shows what that NVSwitch power budget buys. Assume the schedule of the shared reference step, with two activation-sized AllReduces per layer in each forward and backward pass, one after the attention block and one after the feed-forward network; the partition and collective schedule are scenario inputs, not a universal Transformer schedule. During the modeled tensor-parallel forward pass, each GPU holds one shard of each weight matrix and computes its portion of the matrix multiplication. For a Transformer layer with hidden dimension 12,288 (typical for a 175B model), one projection maps each token's internal vector, already split across the 8 GPUs, back to the full hidden width. Each GPU holds a $12{,}288 \times 1{,}536$ slice of that projection's weight matrix, which keeps the full output width but covers only its own eighth of the input. Every output value is a sum over all input values, so one GPU can compute only part of that sum. Each GPU therefore produces a full-width vector of partial sums for every token, and an AllReduce operation sums those partial results across all 8 GPUs over NVLink.

The logical output tensor for one AllReduce contains $1 \times 2{,}048 \times 12{,}288 \times 2$ bytes (FP16). For a microbatch of one 2,048-token sequence, that payload is roughly 50 MB. An ideal ring AllReduce with 8 participants moves $2(N-1)/N = 1.75$ times the payload per GPU (chapter 5 derives this factor), or about 88 MB. At 450 GB/s of one-direction bandwidth, the transfer-time lower bound is approximately 0.196 ms.

With two AllReduces in each of the 96 layers in both the forward and the backward pass, the 384 transfers sum to about 75.2 ms per microbatch before overlap. NVSwitch provides the paths needed

[5] **Crossbar switch**: A topology providing a dedicated path between every input-output pair simultaneously. An $k \times k$ crossbar needs k^2 crosspoints, so its cost grows quadratically with port count; that growth is why a full crossbar stays inside a node and larger systems switch to multi-stage fabrics. Within that bound, the crossbar buys the all-to-all bandwidth that tensor parallelism demands.

↪ How tensor parallelism splits weight matrices across devices is developed in section 6.5.3.

for the collective, but this value remains an ideal bandwidth-only lower bound that excludes startup latency, software overhead, and competing traffic.

Judging that bound needs a compute time to compare it with. At about $6 \times P$ floating-point operations per training token, the group needs 604 ms of computation per microbatch when each GPU sustains 45 percent of its 989 TFLOP/s dense FP16 peak; the fractions below scale with that time, so a longer compute time makes the same communication matter less. Against it, the unoverlapped NVLink lower bound leaves computation at 88.9 percent of combined compute-plus-communication time. The reference job splits the layers across 8 pipeline stages, so each stage's group carries the same fraction of both the transfers and the computation, and this share does not change. Summed over the 96 microbatches of one optimizer step, one stage's transfers come to 0.90 s, the tensor-parallel communication time modeled for the shared reference step. Actual tensor-parallel efficiency can be higher when collectives overlap useful computation or lower when latency, contention, load imbalance, and nonparallel work add delay, so it must be measured rather than inferred from peak link bandwidth alone.

↪ The α-β model, which adds the fixed per-message startup latency this bound omits to the bytes-over-bandwidth transfer time, is developed in section D.1.

Example 2.3: The NVLink bandwidth surprise

Scenario: A team evaluates tensor parallelism over a hypothetical PCIe Gen5 path with 64 GB/s bandwidth, using the payload and the 604 ms per-microbatch compute time derived with the NVLink estimate above.

Diagnosis: The modeled tensor-parallel schedule executes 384 AllReduce events per microbatch across 96 layers. Its PCIe bandwidth-only transfer times sum to 528 ms per microbatch. Without overlap, communication accounts for 46.7 percent of the combined compute and transfer estimate; this is not a measured stall fraction.

Systems lesson: Under the same assumptions, NVLink lowers the transfer bound to 0.196 ms per AllReduce and raises computation to 88.9 percent of the combined time. The comparison explains why this partition favors a fast local fabric. Real collective performance must also include startup, topology, contention, and overlap.

The DGX H100 pairs each GPU with its own InfiniBand ConnectX-7 HCA, which reads and writes GPU memory directly over PCIe through GPUDirect RDMA, so inter-node gradient exchange avoids staging copies in host DRAM while still consuming PCIe and network bandwidth. The host CPU prepares and loads batches, submits kernels, and coordinates the training framework. Sufficient CPU capacity and I/O bandwidth let that work overlap the accelerator fabric's computation and communication; otherwise, slow preprocessing or submission can starve the GPUs even when their local memory and interconnect are fast enough.

Host CPUs also pay a "data center tax" of shared infrastructure work. A profile of Google's fleet found that such common building blocks consumed nearly 30 percent of cycles across jobs (Kanev et al. 2015), and network protocol processing, storage virtualization, and security functions add to that load on cloud hosts. ML nodes can therefore incorporate data processing units (DPUs) or SmartNICs. Devices like the NVIDIA BlueField DPU offload these infrastructure tasks to dedicated ARM cores and hardware accelerators integrated into the network adapter itself. Moving the control plane for RDMA, firewall rules, and storage protocols to the DPU recovers cycles for the ML pipeline: data loading, tokenization, and kernel orchestration on the host CPU. The DPU also acts as an isolated security domain, allowing cloud providers to maintain control over the network and storage layer via the DPU while granting customers full access to the host CPU and GPUs. For our 175B model training, DPU offload can keep RDMA protocol processing outside the host CPU critical path for gradient synchronization, which flows from GPU HBM toward the network fabric without host CPU involvement.

↪ The RDMA and GPUDirect data path that bypasses host memory is developed in section 3.4.1.

2.5.2.1 Alternative node architectures

Spending board area, power, and cost on four NVSwitch chips makes sense when tensor-parallel collectives sit inside every layer, but workloads dominated by embedding capacity, data parallelism, or cost-sensitive serving hit different limits first. Alternative node designs reallocate that budget. An AMD MI300X server emphasizes large HBM capacity on each package, Google TPU pods wire chips

directly into a mesh without a central baseboard switch, and Intel/Habana Gaudi nodes integrate Ethernet RDMA ports directly onto the accelerator die. Comparing them against a DGX H100 node reveals where each architecture spends its scarce budget.

In a switch-rich node like the DGX H100, universal any-to-any local exchange is the priority because each Transformer layer places intra-node collectives inside the model's inner loop. The 900 GB/s H100 NVLink fabric therefore buys a software abstraction, in which the training framework can treat eight accelerators as one tightly coupled tensor-parallel group for the layers that must be split across chips.

Other systems move that boundary. TPU pods use the inter-chip interconnect and a mesh-oriented topology instead of making a central switch the defining component of the node; a TPU v5p chip's vendor-specified ICI bandwidth is 1,200 GB/s. The mesh can reduce switch cost and expose a regular structure to the compiler, but it also makes placement more visible, because the runtime must respect neighbor relationships rather than assuming that every chip pair has the same communication path. That is a reasonable trade when the workload is regular enough for the compiler and collective libraries to map communication onto the topology.

The MI300X shifts the pressure point from local switching toward memory capacity. Its package exposes 192 GB of HBM at 5.3 TB/s, compared with 80 GB on an H100-class accelerator. An eight-accelerator MI300X node would therefore provide 1,536 GB of aggregate HBM, enough to hold the 350 GB FP16 weight tensor for our 175B running model. That does not eliminate the optimizer, gradient, activation, and fragmentation problem that section 2.5.3 quantifies, but it changes how often the system must spill state into slower memory tiers and how much microbatch headroom remains before offload becomes mandatory.

Gaudi takes a different route by integrating RDMA-capable Ethernet into the accelerator and using a more network-centric design for communication. A Gaudi 2 accelerator's 24 integrated 100 GbE ports supply 300 GB/s of aggregate RDMA over Converged Ethernet (RoCE) bandwidth, against the 900 GB/s of local exchange per GPU in an H100-class NVLink node. The newer Gaudi 3 keeps that Ethernet-centric design while carrying 128 GB of HBM per accelerator, so capacity is not the trade being made. The question is whether simplifying the fabric and using Ethernet consistently across node and pod is worth giving up the high local exchange bandwidth of a switch-rich NVLink design. That trade can be attractive when data parallelism, inference serving, or cluster cost dominates, and less attractive when tensor-parallel collectives sit inside every layer.

These alternatives are not a ranking of vendors. They are different placements of scarce resources: switching bandwidth in the node, regular topology in the pod, HBM capacity in the package, or network attachment on the accelerator. For the 175B model used throughout this chapter, the tensor-parallel portion of training makes local accelerator communication unusually important, so the DGX-style answer is natural. In other regimes, especially memory-heavy recommenders, topology-regular training, or cost-sensitive serving fleets, a different boundary can become the one that matters most. Software ecosystem maturity, supply availability, and total cost of ownership then decide whether the theoretical architectural fit becomes a deployable system.

2.5.3 Node memory partitioning

The node's aggregate HBM must hold far more than the weights, and how that state divides across the node's memory tiers decides which parallelism strategies are possible. Consider the memory budget for training our 175B model. The model weights in FP16 require 350 GB. The Adam optimizer maintains two additional values for every parameter (first and second moments) in FP32 precision, adding $175 \times 10^9 \times 2 \times 4$ bytes = 1,400 GB. Gradients require another 350 GB. This subtotal is approximately 2.1 TB, far exceeding both a single accelerator's 80 GB HBM and an 8-H100 node's 640 GB aggregate HBM, and it still omits two terms. Activations are the intermediate results the backward pass needs. The FP32 master weights are a full-precision copy of the parameters that the optimizer updates, kept because an update far smaller than a weight would be lost to FP16 rounding; this is the mixed-precision copy introduced with the precision formats in section 2.4.

A common misconception is that a model's memory footprint equals its weight storage. With those two terms included, the FP32 master weights, optimizer moments, gradients, and activations collectively consume about 7–8$\times$ the memory of the FP16 parameters alone. These components

reside in different tiers of the node's memory hierarchy. For our 175B model, training memory partitions across six distinct components (table 2.11).

Table 2.11: Training Memory Breakdown for a 175B-Parameter Frontier Model: The optimizer state alone (first and second moments in FP32) consumes 4× the memory of the model weights, making optimizer memory the dominant component of training memory. This is why memory optimization techniques focus heavily on optimizer state sharding and offloading.

Component	Precision	Memory per parameter	Total for 175B model
Model weights	FP16	2 bytes	350 GB
Gradients	FP16	2 bytes	350 GB
Master weights	FP32	4 bytes	700 GB
Optimizer (Adam m)	FP32	4 bytes	700 GB
Optimizer (Adam v)	FP32	4 bytes	700 GB
Activations	Mixed	Variable	100–400 GB
Total			**2.9–3.2 TB**

The 2.9–3.2 TB budget of table 2.11 extends the earlier subtotal with the master weights and activations. Its five persistent rows sum to 16 bytes per parameter, or 2.8 TB, the training state given for Archetype A in section 1.6.1.1, and activations add the rest. The Adam optimizer state in FP32, not the FP16 weights, is the dominant term, and the activation peak, which scales with batch size and sequence length, is the most variable. The lesson for parallelism selection is that the FP16 weights are among the smaller pieces of what the node must hold, so a strategy that only shards weights leaves the largest consumers, the FP32 optimizer moments and master weights, untouched.

A DGX H100 node provides 640 GB of aggregate HBM across its 8 GPUs, plus 2 TB of host DDR5 DRAM. The HBM capacity alone is insufficient for the full training state, as the memory arithmetic for each parallelism strategy shows. With pure data parallelism, each GPU must hold replicated state, and even the 2,100 GB subtotal of weights, gradients, and optimizer moments, before master weights and activations, is physically impossible on an 80 GB device. Even ZeRO stage 3, the form of ZeRO sharding that divides all three components among the workers, yields 2,100 GB / 8 ≈ 262.5 GB per GPU when confined to the 8 GPUs of a single node, still more than triple the available HBM.

Fitting the state into HBM therefore requires expanding across multiple nodes so that a larger pool of accelerators shares the memory footprint. The reference job does so with two partitions that follow the bandwidth hierarchy. Tensor parallelism inside each node splits the tensors and operations within every layer across its 8 GPUs, and pipeline parallelism across nodes gives each of 8 nodes a contiguous block of 12 of the model's 96 layers. One model replica therefore spans 64 GPUs on 8 nodes. Splitting every component over those 64 GPUs leaves each one 5.47 GB of FP16 weights, the same again of gradients, and 32.8 GB of FP32 master weights and optimizer moments.

Per GPU, the persistent training state comes to 2,800 GB / 64 ≈ 43.8 GB, which fits on an 80 GB device (85.9 GB in decimal bytes) with about 42 GB left for activations, communication buffers, and workspace. That fit sets the pipeline depth. Four stages would leave 87.5 GB per GPU, more than the device holds, so eight is the shallowest power-of-two pipeline that fits. The 16-way data parallelism across the rest of the cluster then repeats the whole 64-GPU partition, and no state needs to be sharded across replicas.

The 43.8 GB per-GPU subtotal is only the starting point, because the activations on top of it grow with microbatch size and sequence length. If activations and workspace do not fit in the remaining headroom, the next step is to shard the optimizer state across the data-parallel replicas, using the memory-efficient data parallelism developed in section 6.3.3. Dividing the FP32 master weights and moments among the 16 replicas would cut the per-GPU state to about 13 GB, and each replica would then update only its own share of the parameters and pass the result to the others. The memory-budget exercise in section 2.5.3.1 examines a single node without this cross-node partitioning and shows why optimizer offload alone is insufficient there. This arithmetic is why tensor parallelism, pipeline stages, optimizer sharding, activation checkpointing, and memory tiering are used together for large-model training. The node is therefore a high-bandwidth compute and communication unit, not a standalone container for the entire training state.

Partitioning alone can still leave too little room for activations, so it is combined with techniques that shrink the live state. Activation checkpointing stores only a subset of intermediate activations for the backward pass and recomputes the rest from the nearest checkpoint, trading extra forward computation for activation memory, a trade that section 6.3.1.2 develops.

CPU offloading stores state in host DDR5 DRAM and may also move optimizer computation to the CPU. The implementation determines which gradients, weights, or optimizer buffers traverse the PCIe Gen5 path at 64 GB/s. Its cost depends on those bytes, host computation, and the time available to overlap transfers; a short GPU update does not establish cheap offload. NVMe offloading extends storage to SSDs, with the illustrative 5–7 GB/s sequential read rate introducing a further bandwidth constraint that prefetching can hide only when consumption leaves enough time.

An offloaded configuration uses the node's memory hierarchy as a tiered system sorted by how often each piece of state is touched. HBM holds the weight shards, activations, and accumulating gradients needed on every layer of the forward and backward passes. Host DDR5 absorbs the optimizer moments and master weights, which are read and written only once per step during the parameter update. NVMe extends capacity for the largest models by holding seldom-accessed shards that can be prefetched within the step budget. A configuration with sufficient distributed HBM can instead retain optimizer shards on accelerators. The placement follows the capacity and transfer-time calculation, rather than a fixed assignment of every state type to one tier.

The training framework's memory manager moves state between these tiers much as a cache controller does, prefetching the next layer's offloaded weights into HBM during the current layer's computation and evicting state that will not be needed soon. Prefetch depth sets the balance, because too shallow a prefetch stalls computation on PCIe transfers and too deep a prefetch crowds out the current step's activations and gradients. Table 2.12 quantifies the gaps the manager must bridge, more than two orders of magnitude in bandwidth between HBM and an NVMe drive. For our 175B model, a capacity plan must therefore choose among additional sharding, offload, and reductions in live state, and the next budget tests whether a proposed combination actually fits within one node.

Table 2.12: Node Memory Hierarchy: Each tier trades capacity for bandwidth. Latencies are seen from the GPU, so host DRAM includes the PCIe crossing, and a GPU's own share of host bandwidth is limited by its PCIe link rather than by the aggregate DDR5 figure. The training framework's memory manager must orchestrate data flow across these tiers to fit models whose total state exceeds the aggregate HBM capacity.

Memory tier	Capacity	Bandwidth	Latency	Role
GPU HBM3	80 GB × 8 = 640 GB	3.35 TB/s per GPU	< 1 µs	Active computation
Host DDR5	2 TB	hundreds of GB/s aggregate	~1 µs over PCIe	Optimizer state
NVMe SSD	8–30 TB	up to 7 GB/s per drive	~100 µs	Overflow, checkpoints

2.5.3.1 Node-level memory budget

Combining all five memory components for our 175B model on a single DGX H100 node in example 2.4 shows why closing the budget requires coordinating sharding, checkpointing, and offload together rather than tuning any one component in isolation.

Example 2.4: Memory budget for 175B training on DGX H100

Scenario: An engineering team calculates the memory budget for training a 175B parameter model using 8-way tensor parallelism on an 8-GPU DGX H100 node (80 GB HBM per GPU).

Diagnosis: Without ZeRO sharding, each GPU's TP-8 slice of the optimizer state alone requires 175 GB per GPU, which is more than the weights and gradients combined (88 GB). A separate FP32 master-weight shard adds 87.5 GB. Together with retained FP16 gradients, the assumed activations, and workspace, these components bring the per-GPU budget to 375 GB, well past the 80 GB HBM limit. That budget already assumes activation checkpointing; its activation peak is a scenario input that must be replaced by measurements for the chosen microbatch and sequence length.

Systems lesson: Sizing training memory requires managing weights, gradients, optimizer states, activations, and buffers holistically, because relieving one component does not necessarily clear the envelope. Offloading the optimizer moments and master weights to host DDR5 removes those FP32 components and still leaves 113 GB per GPU, because 8-way tensor parallelism keeps a full weight and gradient shard resident on every GPU. Further reductions require offloading or shrinking another live component. Alternatively, adding devices creates an additional data-parallel group across which ZeRO can shard state. The same eight GPUs cannot supply that group, because together they are already the tensor-parallel group of one replica. With no pipeline stages, the tensor-parallel degree times the data-parallel degree equals the device count, so eight-way tensor parallelism on eight GPUs leaves a data-parallel degree of one.

2.5.4 Data loading and I/O pipeline

Even after the node's memory budget closes, its accelerators still need a continuous stream of training data from outside the chassis. If the data loading pipeline cannot keep pace with the accelerators' consumption rate, the most expensive component in the node sits idle; this is the I/O bottleneck. For pre-tokenized language-model training, the raw demand is small. At 500,000 tokens/s stored as 4-byte token IDs, ingestion is only 2 MB/s. The pressure falls instead on the host CPU cores that read, decode, shuffle, and batch examples, on the PCIe paths that data shares with gradient and checkpoint traffic (section 2.5.1.1), and on shared storage that every node reads at once.

GPUDirect Storage can remove the host-memory bounce copy for data that needs no host-side preprocessing, and staging data on local NVMe can absorb synchronized bursts that would otherwise hit a shared filesystem. For node design, the requirement is end-to-end sustained bandwidth at the intended concurrency. Storage servers, network paths, host preprocessing, host memory, and accelerator transfer must all keep pace, because a local benchmark that satisfies one node says little about a shared path serving the fleet. Keeping that node productive over a multi-week run then turns on how often a component fault inside the chassis interrupts execution.

↪ Sizing the input pipeline end to end, with every stage from storage read to host-to-device transfer held at or above the required bandwidth, is developed in section 4.4.

2.5.5 Node health and reliability

When a single GPU suffers an uncorrectable HBM error or a rack power shelf trips, it does not fail in isolation; it takes its entire eight-GPU tensor-parallel group or multi-node rack out of the synchronous training step. A **failure domain** is the set of resources that a single fault takes out of service together, and the physical hierarchy nests these domains inside one another. A failed GPU breaks its tensor-parallel group, which usually spans its node. The node sits in a rack whose shared power shelf, top-of-rack switch, and cooling loop can take every node in it down at once. Racks share row-level power and coolant distribution, and halls share a building's utility feed, cooling plant, and fire compartment.

For a synchronous training job the node is the smallest domain that matters. An eight-way tensor-parallel group within a node becomes unusable when any one GPU fails, because its collectives require all eight participants. The node's effective MTBF depends on the combined failure exposure of every component whose loss stops the job, as well as shared power and cooling failures. The node's figure is therefore an MTBF, the expected interval between successive failures of a unit that keeps running through replacements, while each GPU contributes an MTTF, the distinction drawn with the reliability wall in section 2.1.

Failure exposure spans memory, interconnects, and the electrical and mechanical systems that support them. An uncorrectable error-correcting code (ECC) fault in an HBM stack forces the accelerator out of service and aborts the active training step. High-speed links degrade as repeated thermal cycles stress connectors and cables, which shows up as link retraining or a fatal communication drop. Power-supply components and the pumps and fans of the cooling path wear out under sustained load, and a loss of power or cooling can take several accelerators down together. The accelerator itself usually dominates the count. In Meta's report on the Llama 3 405B pretraining run, faulty GPUs and HBM3 memory errors together caused nearly half of the unexpected interruptions, ahead of software bugs and network switches and cables (Dubey et al. 2024). The ordering in a given fleet

still depends on its equipment and operating conditions, so local incident data should confirm it before it guides spare capacity or preventive maintenance.

Health monitoring can flag a node for investigation before another long job lands on it. A rising rate of corrected ECC errors, rising junction temperature (the temperature at the transistors) at constant load, and intermittent link retraining are the most useful early signals, although no single signal establishes a cause or a time to failure. The fleet telemetry that collects them is developed in section 12.6.1.

↪ Layered fault detection, from ECC and error-rate monitoring to known-answer stress tests on nodes pulled from production, is developed in section 7.4.3.

For our 175B model training across 128 nodes, a composite node-level MTBF assumption of 1,000 hours, which includes GPUs, host components, power, cooling, and network links, predicts about 43 node failures over a run of 336 hours, roughly 3.1 failures/day. The estimate multiplies the node count by the run length and divides by the composite MTBF. A GPU-only bound from the per-GPU MTTF of 50,000 hours would predict only 6.9. The node, not the GPU, is therefore the unit that spare capacity and recovery tooling must be sized for. The same accounting, applied to every component of a facility and then to every GPU of a fleet, reappears at larger scope in the facility reliability controls of the rack and the hours-scale fleet failure interval of section 2.7.1.

For our 175B model, a single DGX H100 node provides 640 GB of aggregate HBM, enough to hold the FP16 model weights but not the full Adam training state without sharding and offload. Training also requires processing hundreds of billions of tokens, and a single node would need more than a year of compute to do so (section 2.7 works through the numbers), so completing training in weeks requires tens to hundreds of nodes. Stacking those nodes into a physical enclosure defines the next level of infrastructure, where the constraints shift from bandwidth and capacity to raw power and heat.

2.6 The Rack

A standard 42U server rack in a traditional data center draws about 12 kW and can be cooled by room-temperature air pushed through perforated floor tiles. The chapter's four-node DGX H100 rack scenario draws 33.5 kW, a figure that includes the rack's allocated cooling overhead. At this density, the engineering constraints shift from silicon and signal integrity to power delivery and thermodynamics. The rack is where the power wall and heat transfer become dominant design forces.

For our 175B model, training across 1,024 GPUs requires 32 racks (4 nodes of 8 GPUs each per rack). Across the cluster, the modeled rack power, which includes each rack's allocated cooling overhead, sums to approximately 1.1 MW. Delivering this power reliably and removing the resulting heat without allowing any component to exceed its thermal limit is a multidisciplinary engineering challenge spanning electrical, mechanical, and civil engineering. A failure at any point in the power delivery chain, from the utility substation to the individual GPU voltage regulator, can halt the entire training run, wasting hours of computation and potentially corrupting the training state.

Definition 2.5: Rack

Rack is the physical infrastructure unit, a standardized 42U enclosure, that houses multiple compute nodes, a top-of-rack (ToR) switch (the first network aggregation point connecting all nodes in the rack to the broader cluster fabric), power distribution units, and cooling distribution manifolds, defining the granularity at which power and cooling capacity must be provisioned.

1. **Significance**: AI rack power density has grown dramatically. The chapter's four-node DGX H100 scenario contains 32 GPUs at 700 W each and draws 33.5 kW once host, network, conversion, and allocated cooling overhead are included. This exceeds the roughly 12 kW typical of general-purpose data center racks and requires purpose-designed electrical distribution and cooling capacity. Whether that cooling is engineered air or liquid depends on the system configuration and facility design.
2. **Distinction**: A node packages compute; a rack integrates shared power, cooling, and network distribution. Components and nodes remain individually serviceable, but their

maintenance and replacement must respect the capacity and failure boundaries of those shared systems.

3. **Common pitfall**: A frequent misconception is that rack placement is a logistical afterthought. Nodes within the same rack share a ToR switch and have single-hop connectivity; nodes in different racks cross at least two switch hops. Placing pipeline stages within a rack can reduce inter-rack traffic when the actual paths and available capacity support the stage-boundary exchanges.

2.6.1 Power and thermal envelopes

Sustaining the compute rate that HBM and Tensor Cores make possible requires a power-delivery and cooling system sized for the workload. Delivery itself adds heat before any computation does, because each of the five voltage-conversion stages between the medium-voltage utility grid and the sub-volt GPU die loses part of the power it passes on (figure 2.12).

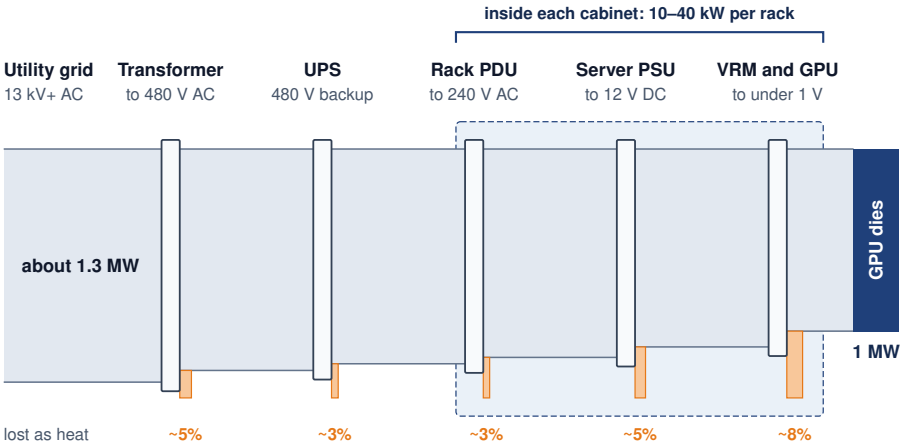

Figure 2.12: The Power Delivery Path: Energy passes through five conversion stages between the medium-voltage grid and the sub-volt supply at the GPU die, and each stage loses a few percent as heat; the band's height is drawn proportional to the power it carries. The per-stage losses shown are illustrative round values, not measurements of a particular facility; compounded, they lose about a fifth of the power, so delivering 1 MW to the dies draws about 1.3 MW before cooling. PUE uses an IT-equipment boundary and includes facility conversion losses as well as cooling; these losses must not be counted twice. The critical engineering challenge is the rack power distribution unit (PDU) to voltage regulator module (VRM) transition, where 10–40 kW of power must be delivered within a single cabinet.

The illustrative conversion losses in figure 2.12 require roughly 1.3 MW upstream to deliver 1 MW to the dies, before heat removal is counted. Facilities summarize that overhead with power usage effectiveness (PUE), total facility power divided by IT-equipment power, and because the IT boundary already contains some of these conversion losses they must not be counted twice. Section 15.4.1.3 defines the ratio and draws its boundary.

Every floating-point operation dissipates energy as heat. Dense Tensor Core activity during large matrix multiplications drives heat proportional to the switching activity of billions of transistors. Fleet design must pair each accelerator's sustained power envelope with any documented transient excursion requirement. At dense fleet scale, this produces the power wall, the point at which facility power delivery and heat removal, rather than available arithmetic units, determine how much compute can be installed and sustained.

The thermodynamic limit (principle 4), which holds that every joule a computation draws ends as heat and that a facility can run only as fast as it removes that heat, makes heat removal part of usable compute capacity. A facility can reach its power or cooling envelope while arithmetic capacity is still available to purchase, so the scarce quantity becomes watts per unit of floor area rather than

FLOP/s. The power density wall (principle 3) adds that spatial constraint, bounding cluster scaling by the watts each rack can dissipate rather than only by silicon area or floor space. If a proposed rack exceeds the engineered airflow's heat-removal capability, the design must lower its density or move to liquid, which at the 100 kW class of high-density training racks is part of the rack design rather than an option.

Definition 2.6: Thermal design power (TDP)

Thermal design power (TDP) is a vendor-defined sustained thermal design target in watts for a specified product configuration and operating condition.

1. **Significance**: The H100 SXM5 configuration used in this chapter carries a 700 W design target. Sustaining the intended operating point requires a thermal path and facility loop qualified for that target under the vendor's inlet and coolant conditions.
2. **Distinction**: TDP does not denote one universal instantaneous-power maximum. Total graphics power (TGP), board power, configurable limits, and transient specifications describe related electrical constraints under conventions that differ by vendor and product.
3. **Common pitfall**: A frequent misconception is that the advertised wattage alone specifies rack power and cooling. Fleet design must combine the sustained thermal target, configured power limit, workload behavior, component overhead, and transient-power requirement.

The trajectory of TDP records a scaling regime reaching its limits. For three decades, Dennard scaling[6] allowed chip designers to shrink transistors while maintaining approximately constant power density. Smaller transistors required lower voltage, so packing more transistors into the same area did not increase heat output per unit area. As voltage scaling and leakage limits ended that regime, architects had to treat energy, cooling, and parallelism as first-order constraints rather than free byproducts of process shrinks (Dennard et al. 1974; Hennessy and Patterson 2019). Since then, each generation of accelerator has had to trade performance against power and heat. The era of "free" performance improvements through process shrinks is over; every gain in throughput now has to be paid for in watts, cooling capacity, or specialization.

The V100 (2017) operated at 300 W TDP, the A100 (2020) increased to 400 W, the H100 (2022) reached 700 W, and the B200 (2024) draws 1000 W. In seven years, TDP has more than tripled, growing at approximately 19 percent per year.

If this growth rate continued as a planning scenario on a two-year product cadence, a generation after Blackwell would approach 1,400 W per accelerator, and the generation after that would reach about 2,000 W. At that level, a single 8-GPU node would consume 16 kW of GPU power alone, requiring liquid cooling infrastructure (section 2.6.2) that can remove heat at rates approaching those of industrial process cooling equipment. The applicable sustained power envelope is not merely a specification on a data sheet; it propagates through every level of the infrastructure stack. The power-delivery chain, cooling system, rack design, and facility electrical infrastructure must accommodate deployed accelerators, their documented transients, and the plausible envelope of future refresh cycles.

The propagation is visible at both ends of the stack in published data. Vendors publish a power rating for each accelerator package, and figure 2.13 (a) collects those ratings across 14 NVIDIA, Google, AMD, and Intel parts. The ratings follow each vendor's own convention (TDP, maximum board power, or maximum TGP), so they bound design envelopes rather than measure consumption, but they climb steadily from one generation to the next. Panel (b) plots the facility power of supercomputers that Epoch AI tracks for AI work, together with the largest one operating at each date (Epoch AI 2026a; Pilz et al. 2025). From 2013 through 2021 that record stayed between 20 MW and 30 MW, a level set by national high-performance computing systems. The Frontier supercomputer raised it to 40 MW in 2022, and AI training clusters then lifted it to about 350 MW by 2025, a 9× increase in three years. Over the same three years the rating of one package rose from the H100's 700 W TDP to the B300's 1,400 W maximum TGP, roughly a 2× increase across the two rating conventions. Hotter packages explain only a small part of the jump. The rest comes from the number

[6] **Dennard scaling**: Named after Dennard et al. (1974) at IBM, who showed how MOSFET dimensions, voltage, and current could scale so that power density stayed roughly constant. Later voltage-scaling limits and leakage effects ended that constant-power-density regime; Hennessy and Patterson (2019) describe the resulting post-Dennard power wall as one reason modern architecture turned toward parallelism and specialization. For ML infrastructure, the V100-to-B200 TDP trajectory (300 W to 1,000 W) is an illustration of the post-Dennard planning problem: cooling and power delivery are now co-equal design constraints with the silicon itself.

of accelerators a single system joins together and the facility built to feed them, which is why the power wall at fleet scale is a provisioning problem for the building as much as a thermal problem for the chip.

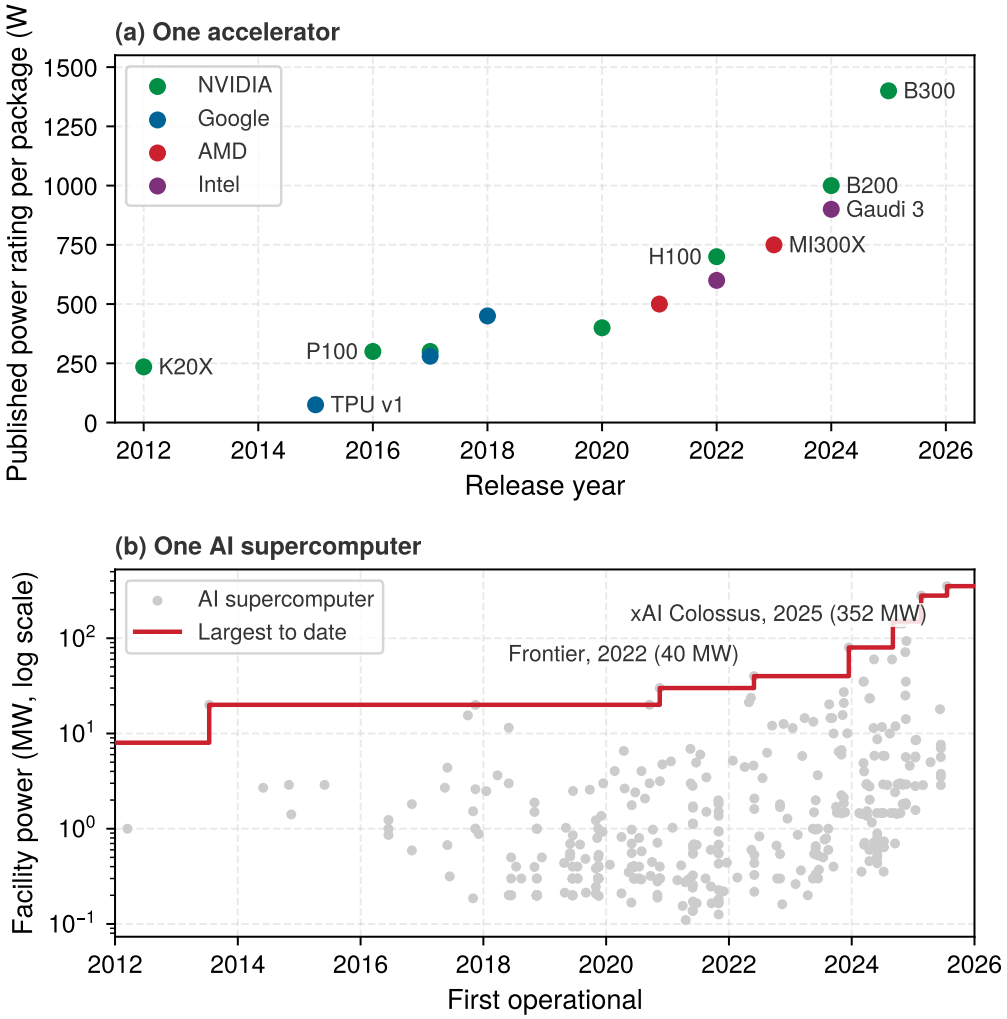

Figure 2.13: Package Power versus System Power: (a) Published power rating of one accelerator package by release year. Ratings follow each vendor's convention (TDP, maximum board power, or maximum TGP) and describe design envelopes, not measured workload power. (b) Facility power of AI supercomputers by first operational date on a logarithmic axis; the red line tracks the largest system operating at each date. The record held near 20 to 30 MW for most of a decade before AI training clusters raised it roughly ninefold between 2022 and 2025. Sources: (a) NVIDIA (2013), Foley and Danskin (2017), Choquette et al. (2018), Choquette et al. (2021), Choquette (2023), NVIDIA (2024), NVIDIA (2025), Jouppi et al. (2021), Advanced Micro Devices (2021), AMD (2023), Intel Habana Labs (2023), Intel (2024); (b) Epoch AI (2026a).

Accelerators implement power management to operate within their TDP envelope. When the workload does not fully use all SMs (for example, during the communication phase of a training step when the GPUs are waiting for AllReduce to complete), the GPU firmware can disable the clock signal to idle circuitry through **clock gating**. Clock gating eliminates dynamic switching power and frees thermal headroom for the active SMs. A separate mechanism, dynamic voltage and frequency scaling (DVFS), adjusts the chip's clock frequency and supply voltage to stay within power and thermal limits. During a large matrix multiplication that activates all SMs simultaneously, DVFS reduces the clock frequency to keep total power within TDP. When some SMs are idle, the active

SMs can boost frequency because total chip power is below TDP, so the actual clock, and therefore the actual throughput, varies throughout a training step.

Dynamic power management means that the actual power draw of a GPU varies continuously during training. A typical training step might see power swing from 400 W (during communication phases when most SMs are idle) to 700 W (during the forward and backward pass when all Tensor Cores are active) and back, with each transition occurring in microseconds. The voltage regulator modules (VRMs) on the baseboard, which convert the board supply to the sub-volt core voltage, must respond to these transitions fast enough to maintain a stable supply voltage despite the rapidly changing current draw.

DVFS also couples software optimization to the power envelope. A kernel that achieves high SM occupancy (many active warps, or scheduled thread groups, per SM) draws more power, potentially triggering a frequency reduction that partially offsets the occupancy benefit. Conversely, a kernel with moderate occupancy may run at a higher clock frequency, achieving comparable throughput at lower power. Power-normalized throughput (TFLOP/s per watt) is therefore a more robust comparison metric than raw TFLOP/s, and it explains why the same GPU can show different sustained clock frequencies depending on the workload and the cooling solution.

At fleet scale, DVFS also affects power planning. Average workload draw can fall below the configured limit because communication phases and incomplete occupancy leave circuitry idle, but averages do not replace sustained and transient design cases. Some operators deliberately lower accelerator power caps to fit more devices within a fixed facility budget, accepting workload-dependent throughput loss in exchange for lower per-device draw. Power capping is a form of the compute-efficiency trade-off at the operational level.

The physical challenge of 700 W lies in its heat flux. An H100 die measures approximately 814 mm^2 (roughly 29×28 mm). Dissipating 700 W from this area produces a heat flux of approximately 86 W/cm^2. For comparison, the surface of an electric stovetop burner at high heat produces approximately 10 W/cm^2. An H100 die generates 8.6× the heat flux of a stove burner, concentrated on an area the size of a large postage stamp.

The B200's 1,000 W envelope spans a dual-die package. In this comparison, the average heat flux remains comparable while the package heat-removal rate increases by 43 percent. Cooling must keep junction temperature and memory temperature within the applicable product limits under the specified inlet and load conditions. A single temperature threshold cannot stand in for those separate requirements.

The thermal interface material (TIM) between the package and its cooling assembly is one part of that heat-transfer path, and it is the part that ages. Training alternates compute-dense phases with communication phases, so the die heats and cools many times per minute; under that cycling, a phase-change or paste TIM can pump out (migrate away from the contact surface) or crack. Increased thermal resistance can raise temperature at the same load, but the performance effect depends on the remaining thermal margin and the device's control policy. A temperature trend alone does not identify the cause; inlet conditions, flow, contact, and workload also require inspection. When thermal limits reduce clocks without stopping execution, the result is a gray failure, a performance degradation that monitoring must distinguish from an application or network bottleneck.

Arithmetic per watt offsets only part of rising TDP, and a larger capability target can increase total work despite improved efficiency, so facility planning must compare the proposed device count, power envelope, and completion time together. Packing accelerators close together shortens communication paths but concentrates heat. A four-node DGX H100 rack already draws 33.5 kW including its allocated cooling overhead, several times the load of a conventional server rack. The accelerators of a 10,000-GPU cluster, the round fleet size this chapter uses for warehouse-scale examples, alone draw 7 MW before hosts, network, and cooling are counted. Rack and pod designs must reconcile that proximity with enough space and infrastructure to deliver power and remove heat across all participating nodes.

Power-envelope implications recur throughout the infrastructure hierarchy. At the node level, configured limits constrain how many accelerators can share a chassis (section 2.5). At the rack level, aggregate sustained load and heat flux help determine whether engineered air or liquid cooling can meet the design point (section 2.6.2). At the pod level, coincident demand and transient behavior shape megawatt-scale power delivery (section 2.7). Every level above the die must therefore manage

both sustained heat and electrical variation. The facility envelope also moves slowly, because upgrading transformers, switchgear, and the cooling plant is a construction project that can take longer than delivering the accelerators. Accelerator selection must compare useful workload throughput within that facility envelope, rather than treating power as a cost added after the compute decision.

2.6.1.1 Power delivery networks and transient dynamics

Fitting redundant AC power supply units inside every eight-GPU chassis consumes enclosure volume needed for cooling manifolds and high-speed interconnects, especially as rack loads climb toward 100 kW in liquid-cooled clusters. Many dense-rack designs therefore replace per-chassis AC supplies with rack-level DC power distribution. In this architecture, a centralized power shelf converts mains AC to a 48V DC busbar that runs the height of the rack. Power reaches the server backplane through the 48V bus, and local DC-to-DC converters supply the lower voltages required by the components. Consolidation can recover chassis volume for cooling loops and high-bandwidth interconnects, but it does not by itself establish fewer conversion stages or a fixed efficiency gain. The comparison must account for each design's actual stages and conversion efficiency across the intended load range. Shared power shelves and busbars also create a rack-level failure boundary whose redundancy and serviceability must be qualified.

A training step makes each accelerator's current draw jump. When a collective barrier releases a node's GPUs from a communication phase into dense arithmetic, each one steps from a few hundred watts toward its full power envelope within microseconds, and the barrier correlates those steps across the node. The power delivery network (PDN) must hold the supply voltage steady through that step, and the difficulty grows as the voltage falls from the 48V busbar to the sub-1V rail at the die.

At a fixed power, current scales inversely with voltage, so the core rail carries far more current than the busbar; the voltage regulators transform voltage and current between those boundaries. What disturbs the supply is the rate of change of that current, its slew rate di/dt. Every conductor loop has inductance, a property that resists changes in current, so a fast current step produces a momentary voltage dip, or droop, even when average power remains within the facility budget.

Definition 2.7: Power delivery network (PDN) and transient dynamics

Power delivery network (PDN) is the electrical distribution hierarchy that supplies the processor through regulators, interconnects, and decoupling capacitors. Parasitic loop inductance contributes to voltage fluctuations during changes in current.

1. **Significance**: Synchronized workload transitions can correlate accelerator demand. Supply excursions outside the device's specified limits can cause timing errors, protective throttling, or resets.
2. **Distinction**: Local electrical transients require a faster response than facility power-control loops can provide. Regulators and nearby charge storage address different timescales within the same hierarchy.
3. **Common pitfall**: A frequent misconception is that the full core-current slew acts on a single inductance summed from rack to die. Regulation and decoupling change the current waveform at each stage; the relevant inductance and current must belong to the same path.

For a particular inductive path, the voltage excursion contributed by its inductance follows equation 2.2:

$$\Delta V = L_{\text{parasitic}} \frac{di}{dt} \tag{2.2}$$

Here $L_{\text{parasitic}}$ is the loop inductance of that path, and di/dt is the current slew through it. Local capacitors initially supply part of a load-current increase, so the current slew seen upstream need not equal the slew at the core. A complete rail model also accounts for resistance, capacitance, and regulator response.

A worked estimate for one accelerator shows why the hierarchy matters. When an H100 moves from its communication-phase draw of about 400 W to its 700 W envelope on a 0.8 V core rail, its core current rises by about 375 A. Suppose the board loop between the voltage regulator and the package has 50 pH of inductance. If package capacitors supply the fast edge so that this loop sees the step over 1 µs, equation 2.2 gives a droop of about 19 mV, or 2.3 percent of the rail. If the same step reached that loop in 100 ns, the droop would be about 190 mV, roughly 23 percent of the rail and far outside a 5 percent operating band. The inputs are scenario assumptions, but the tenfold sensitivity to rise time is physical. The same current step stays inside tolerance or corrupts timing depending on which charge store answers it first.

The collective barrier that aligns transitions from communication to arithmetic raises the correlation of demand across a rack, but it does not imply that every accelerator switches on the same nanosecond or that their core currents flow through one shared low-voltage loop. Engineers must evaluate both each device's local voltage excursion and the aggregate demand seen by upstream supplies. A rising load can produce voltage droop, where droop below the specified operating minimum can invalidate timing margins, while an abrupt load reduction can produce overshoot that must remain within the device's specified limits as well. Neither the magnitude nor the resulting failure mode follows from workload synchronization alone. The mitigation is a hierarchy of charge storage and control, in which each tier answers the transients too fast for the tier above it (table 2.13), with placement and parasitic impedance determining which frequencies each tier can address.

Table 2.13: Decoupling Tiers in the Power Delivery Network: Each tier supplies charge on a timescale the slower tier above it cannot reach, so a fast edge that bypasses a tier lands on a larger inductance.

Tier	Charge store or control	Response timescale	Transients it absorbs
On-die	On-die capacitance	Sub-nanosecond	Switching activity within a kernel
Package	Package and interposer capacitors	Nanoseconds	Kernel launches and streaming-multiprocessor activation edges
Board	Bulk capacitors, regulator control loop	Microseconds	Communication-to-compute phase changes of one accelerator
Facility	Uninterruptible power supply (UPS), utility	Milliseconds to seconds	Synchronized power swings across racks and the cluster

The facility row of table 2.13 is where the barrier's correlation appears in aggregate. For an illustrative calculation, assume 1,024 GPUs change together from 400 W to 700 W each over 100 µs. The assumed aggregate step is 307 kW, giving a ramp rate of 3.07 GW/s. This is a load-side scenario, not a measured utility input. No single regulator carries this step, because each accelerator draws its share through its own board, yet the shared upstream supplies see the sum.

Workload-level controls complement these hardware tiers. Power ramping can reduce abrupt demand changes where the hardware and firmware support it, but a design must verify those controls rather than assume that every kernel launch receives automatic smoothing. The fleet-level lesson is that average power capacity and transient voltage integrity require separate checks.

↪ How energy buffering presents this ramp to the grid as a smoother, more easily connected load is developed in section 15.6.2.

2.6.2 Cooling class and heat removal

The power path ends in heat. Under the thermodynamic limit (principle 4), nearly every joule an accelerator draws leaves as heat through switching, leakage, and resistive losses. During a transient some of that energy warms the chip and its surroundings, but in steady operation the stored thermal energy stops rising, so cooling must remove heat at about the rate the rack consumes electricity. When removal falls behind generation, temperature climbs through the available thermal margin until throttling lowers clocks and throughput, and sustained overheating can damage the hardware. The rack's cooling class, air or liquid, is the choice of fluid that carries this heat away, and it bounds how dense the rack can be.

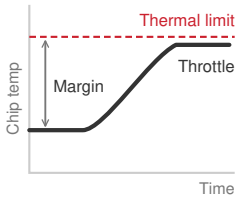

When cooling falls behind, temperature climbs through the thermal margin until throttling caps it.

Heat transfer explains why density reshapes the rack. Air carries far less heat per unit volume than liquid, so an air-cooled rack needs enough airflow, pressure, and temperature rise to carry away its heat. How much a hall can supply depends on the server and facility design rather than on one rack-power threshold.

Air cooling, the dominant technology for decades, works by blowing room-temperature air across heat sinks attached to the chips. The air absorbs heat at a rate determined by its specific heat capacity, roughly 1.0 kJ/kg/K. The heated air is exhausted from the rear of the rack, typically 15–20 degrees Celsius warmer than the inlet, and directed to a computer room air conditioning (CRAC) unit that cools it before recirculating.

Air is a poor thermal conductor. Its thermal conductivity is only 0.026 W/m/K, compared to 0.6 W/m/K for water and 400 W/m/K for copper. Removing a large heat load with air requires substantial volumetric flow. Fan power depends on that flow, pressure drop, and fan efficiency; rack power alone does not specify a fixed cooling fraction.

The heat that a coolant stream removes depends on how much fluid flows, how much heat each kilogram absorbs per degree of warming, and how far the fluid warms (equation 2.3):

$$P_{\text{thermal}} = \dot{m} \times c_p \times \Delta T \tag{2.3}$$

where P_{thermal} is the heat removed (watts), $\dot{m}$ is the mass flow rate (kg/s), c_p is the specific heat capacity (J/kg/K), and ΔT is the temperature difference between outlet and inlet. For air with c_p of 1,005 J/kg/K and a typical ΔT of 15 K (inlet at 20 degrees C, outlet at 35 degrees C), removing 100 kW requires a mass flow rate of $100{,}000 \div (1{,}005 \times 15) \approx 6.6$ kg/s. At sea-level air density of 1.2 kg/m^3, this corresponds to a volumetric flow rate of 5.5 m^3/s, or approximately 11,713 CFM (cubic feet per minute). Moving this much air through the confined space of a server rack requires powerful fans that themselves consume substantial power.

At higher power densities, the airflow and pressure required to cool a rack raise fan energy and can make air cooling impractical before the IT load itself reaches a power limit. The exact crossover depends on rack geometry, inlet temperature, pressure drop, and fan-control policy rather than on one universal rack-power threshold.

Water has a specific heat capacity of 4.18 kJ/kg/K, over four times that of air, and a thermal conductivity roughly 23× higher. Its much greater density also lets a small fluid volume carry substantial heat. The difference appears in how much fluid must flow to remove 700 W from a single GPU. Air at a 15-degree temperature rise requires tens of liters per second of airflow (a small wind tunnel). Water at the same temperature rise requires only about 0.01 liters per second (a thin stream). This thousands-fold difference in volumetric flow permits compact liquid paths, but it does not determine pumping energy. Pressure drop through cold plates, pipes, and fittings, together with pump efficiency, determines the power needed to sustain that flow.

At rack scale the same equation sizes the liquid loop. Removing the modeled 120 kW load of a GB200 NVL72 rack through water that warms by 10 °C takes about 170 liters per minute, and a smaller temperature rise needs proportionally more flow. A real design must use the heat fraction the liquid loop actually captures and account for residual air cooling; flow capacity, pressure, leak containment, and loss-of-flow response are separate requirements, which the facility controls of section 2.6.4 treat as failure domains.

Direct-to-chip liquid cooling routes coolant through cold plates mounted on GPU packages. The cold plate contacts the heat spreader through thermal interface material; the resulting thermal resistance depends on the package, interface, plate design, and flow. The coolant carries heat through manifolds and pipes to a coolant distribution unit (CDU) at the rack or row level. The permissible supply temperature depends on the equipment, so liquid cooling need not require chilled water.

In a two-loop design, the CDU transfers heat from the chip-level coolant to the facility water loop, which carries it toward cooling towers or dry coolers. Isolation lets each loop maintain its required fluid chemistry and pressure. The chip-level coolant must satisfy the equipment's compatibility and cleanliness requirements to control deposits, corrosion, and blockage; purity alone does not establish suitability. Liquid cooling can reduce server airflow requirements, although memory modules (DIMMs), voltage regulator modules (VRMs), and other components without cold plates may still need fans, and lower fan demand can also reduce noise.

A more aggressive approach, immersion cooling, submerges entire server boards in a tank of nonconductive dielectric fluid. The fluid absorbs heat from wetted surfaces, replacing server airflow; high-heat-flux components may still need heat sinks or directed flow, and flow, local heat flux, material compatibility, and heat-exchanger capacity still determine component temperatures. Single-phase immersion cooling uses a fluid that remains liquid throughout the process, with heat carried

away by convection currents in the tank. The heated fluid rises to the surface, is pumped through a heat exchanger to reject the heat to the building's chilled water loop, and returns to the bottom of the tank. Two-phase immersion cooling takes this further. Fluid boils at heated surfaces, carries latent heat to a condenser, and returns as liquid, with heat removal depending on the fluid, pressure, condenser capacity, and operating conditions. The complete facility still spends energy on heat rejection and electrical distribution, so the cooling mechanism alone does not settle facility energy.

The trade-off is serviceability. Accessing a failed component requires handling fluid-wetted hardware, and the dielectric fluids themselves are expensive ($20–$50 per liter). A single immersion tank holding four server boards may contain 500–1,000 liters of fluid, representing $10,000–$50,000 in coolant cost alone. The operational procedures for immersion-cooled facilities differ sharply from air-cooled ones, requiring specialized training for technicians and different approaches to cable management, since all wetted materials must be compatible with prolonged fluid exposure. Depending on the fluid and its additives, metals can corrode, while polymer insulation and seals can swell or degrade. Qualifying compatible materials adds engineering work and can constrain component choices. These serviceability costs enter the design decision alongside the heat paths, binding limits, and parasitic loads of the air and liquid regimes (table 2.14).

Table 2.14: Cooling Architecture Trade-Offs: Feasible rack density follows from the particular equipment and facility design. The example design points differ by 10× in rack power, yet there is no universal rack-power threshold or PUE that separates air and liquid cooling. Compare qualified operating envelopes and total facility energy.

Design question	Air cooling	Direct-to-chip liquid cooling
Heat path	Chip to heat sink to moving air	Chip to cold plate to coolant loop
Binding limits	Airflow, pressure drop, inlet temperature, and allowable component temperature	Cold-plate resistance, flow, coolant temperature, and heat-rejection capacity
Parasitic loads	Server and facility fans; conditioning where required	Pumps, CDU, heat rejection, and any residual air cooling
Service constraints	Air containment, filters, and access paths	Fluid compatibility, leak management, and maintainable connections
Facility metric	Measure complete IT and non-IT energy	Measure the same boundary; liquid cooling does not determine PUE by itself
Example design point	Enterprise rack near 12 kW; best air-cooled facility profile PUE 1.12	GB200 NVL72 rack at 120 kW; liquid-cooled facility profile PUE 1.06

Direct-to-chip liquid cooling supports higher heat flux than forced air because cold plates route chip heat directly into circulating fluid with high volumetric heat capacity before reaching facility heat rejection, whereas air systems rely on convective dissipation across finned heat sinks (figure 2.14). Ambient wet-bulb temperature (the lowest temperature that evaporation can bring outdoor air to, which bounds evaporative heat rejection) and rack power density determine whether an air, liquid, or hybrid loop satisfies the chip junction temperature envelope.

Two anchors bound the choice for this chapter's racks. The four-node DGX H100 scenario at 33.5 kW can use engineered air cooling in a suitable facility, while racks in the 100 kW class of high-density training racks need liquid, which raises heat-removal capacity and can reduce fan energy. Between the two anchors, the facility's airflow capacity decides. Cooling class is therefore a configuration constraint. Rack layout and heat-removal capacity, not accelerator count alone, determine whether the intended configuration can sustain its workload or must reduce density or power.

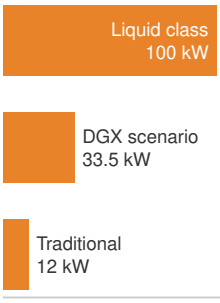

Engineered air can still cool the DGX rack; the 100 kW class needs liquid.

↪ What heat removal adds to facility energy (PUE), and the water that heat rejection consumes, are developed in section 15.6.3.

2.6.3 Rack design considerations

Packing 33.5 kW of compute and thermal hardware into a standard 42U frame leaves almost no room for conventional server layouts. In an enterprise 42U rack, servers are installed as independent 1U or 2U pizza-box units, each with its own fans, power supplies, and cable connections. A single DGX H100 instead occupies 8U and contains 8 GPUs, the NVSwitch fabric, multiple power supplies,

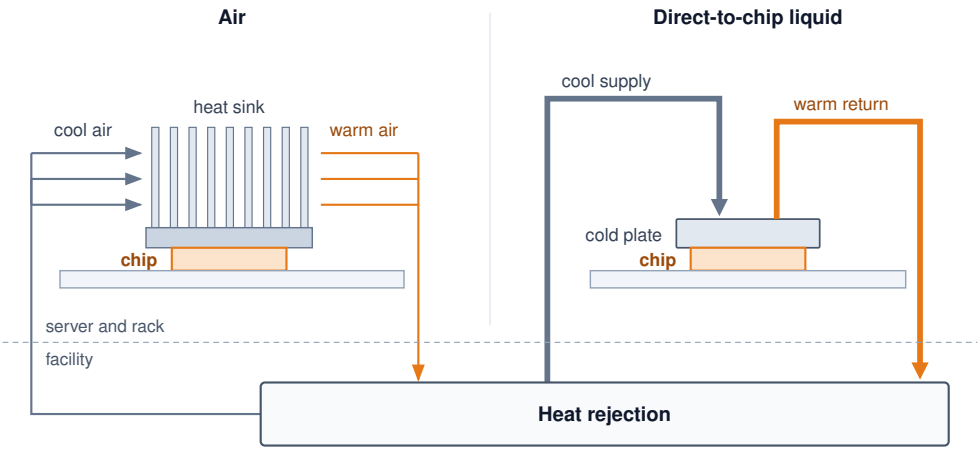

Figure 2.14: Air and Liquid Heat Paths: Air transfers chip heat through a heat sink and moving air; direct-to-chip liquid cooling transfers it through a cold plate and coolant. Both require heat rejection. Capacity, parasitic energy, and water use depend on the engineered system rather than the cooling medium alone.

and (in liquid-cooled configurations) coolant manifolds with quick-connect fittings. Four DGX units fill 32U of a 42U rack, leaving only 10U for networking switches, cable management, and PDUs.

The transition from traditional server form factors to ML-optimized designs is driven by the same power density challenge that motivates engineered airflow or liquid cooling. A traditional 1U server dissipates 300–500 W and requires only modest airflow for cooling. A DGX H100 node dissipates approximately 8.4 kW across its 8 GPUs, NVSwitches, and supporting components, requiring either high-volume forced airflow (in air-cooled configurations) or liquid cooling plumbing (in liquid-cooled configurations). The form factor must accommodate both the compute hardware and the thermal management infrastructure, which is why ML nodes are substantially taller (8U vs. 1U) than traditional servers.

Cable management is a nontrivial engineering challenge. Each DGX H100 node has eight InfiniBand cables (one per GPU), two Ethernet management cables, power cables, and, for liquid-cooled units, coolant hoses. A fully populated rack contains over 40 high-speed cables whose layout is governed by three strict physical constraints: airflow clearance, minimum bend radius, and electromagnetic interference. In air-cooled designs, cables must not obstruct intake or exhaust airflow, because even a localized blockage creates hot spots that trigger thermal throttling. High-speed cables simultaneously enforce a minimum bend radius (for installed fiber optic cable, typically 10–15$\times$ the outer diameter, per the manufacturer's specification (The Fiber Optic Association 2020)), below which optical fibers suffer attenuation and induce bit errors or complete link drops. Densely bundling high-speed copper cables can also induce electromagnetic crosstalk that degrades signal quality. In large installations, the cable plant alone can take weeks to install and test, and cable routing errors are one of the most common causes of post-installation debugging delays.

The rack's network attachment needs the same density-driven rethinking, starting with the top-of-rack (ToR) switch. In traditional data centers, a ToR switch provides 1–10 Gb/s Ethernet connectivity to each server, aggregated into uplinks to higher-level switches. In an ML rack, the ToR switch is an InfiniBand or high-speed Ethernet switch providing 400 Gb/s per port, and it must handle the bursty, synchronized traffic patterns of distributed training. In a conventional design all eight GPUs of a node share one ToR switch, so a data-parallel AllReduce converges on that switch and its uplinks. A rail-optimized design instead gives each GPU index its own leaf switch, the first switch tier that endpoints attach to (GPU 0 of every node on leaf 0, GPU 1 on leaf 1, and so on), so same-index collectives stay on one switch. The price is extra switches and cabling, and the benefit depends on how the job's process groups map onto these per-index rails.

↪ How rail-optimized topologies align cabling with collective traffic is developed in section 3.5.3.

2.6.4 Facility reliability and environmental controls

A coolant leak or electrical fault in a single DGX H100 rack threatens $1.4 million in hardware and can halt a multi-rack training job across the upper levels of the failure-domain hierarchy introduced in section 2.5.5. Data centers housing ML infrastructure therefore pair physical security with continuous environmental monitoring that becomes part of the compute design rather than a facilities afterthought. Beneath raised floor tiles and along liquid cooling piping, water sensors identify leaks within seconds and close automatic coolant isolation valves, while humidity controls keep the air within a moderate band, dry enough to prevent condensation and humid enough to limit static discharge. Inside the cabinets, aspirating smoke detectors sample rack air continuously to catch combustion byproducts at concentrations far below what sprinkler-activating detectors can sense. Vibration sensors on the building structure and rack frames flag seismic events or construction shocks before they can damage disk drives or loosen cable connections.

Cooling failures threaten both hardware and useful fleet work. When coolant flow stops, a dense rack loses heat-removal capacity and must throttle or shut down before thermal runaway. N+1 redundancy in CDUs (one spare beyond the number the load requires), temperature monitoring, and automatic load shedding limit the affected capacity. An emergency air path can support reduced-power operation where the facility design provides one.

Serviceability also changes the failure budget. Leaks can damage equipment and drain cooling capacity, which is why the leak sensors above drive isolation valves, while quick-disconnect fittings allow node replacement without draining the entire loop. Their seals and coolant compatibility therefore belong in the maintenance plan. The fleet-level consequence depends on how much synchronized work shares the failed cooling domain.

The economics of cooling reliability shift when moving from independent inference servers to tightly coupled training clusters. In a distributed training run using synchronous parallelism, a cooling-loop failure that idles one slice of the cluster can halt the entire job. Consider a cooling failure that triggers a thermal shutdown of the 256 GPUs served by one coolant distribution unit within a 10,000-GPU cluster, about ten times the 1,024-GPU reference cluster. The example counts lost compute capacity, excluding repair and replacement costs. If the repair time for a CDU pump is 4 hours, the immediate loss of 256 GPUs at the fleet's reference in-service rate of $2 per GPU-hour is $2,048. Assuming the job cannot continue elastically and the remaining 9,744 GPUs remain unavailable for other work, the synchronous blast radius inflates the cost to $80,000 in lost compute time. Adding checkpoint retrieval and rollback of 30 to 60 minutes of lost computation brings the illustrative lost-compute total to $90,000–$100,000 before any additional recovery delay. The energy the rollback repeats belongs to the facility energy account of section 15.6.3. Whether N+1 cooling redundancy is worth its price follows from comparing these expected outage losses and recovery behavior with that price, rather than from treating redundancy as an unconditional requirement.

Coolant condition is another reliability input. **Deionized water** and compatible inhibitors control **galvanic corrosion**, the damage caused by electrochemical differences between metals in the loop. **Biofilm** on heat-transfer surfaces impedes cooling and increases pumping demand. Maintenance must preserve coolant chemistry and heat-transfer performance, because a fouled or corroding loop loses capacity before it fails outright.

Fire containment and power isolation bear directly on reliability and uptime. A fire in a single rack can shut down an entire data center hall for days while the fire suppression system is recharged and the affected area is decontaminated. The 2021 OVHcloud fire in Strasbourg shows how a single-building incident, absent the right compartmentalization, escalates into a site-wide failure domain.

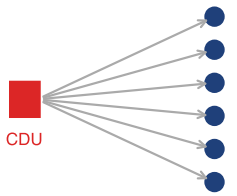

One cooling failure can idle the whole training job.

War Story 2.1: The data center that became the failure domain (2021)

Context: OVHcloud operated four data centers at its Strasbourg site, serving customers whose systems depended on the site's physical power, fire, and recovery design (OVHcloud 2021).

Mechanism: On March 10, 2021, a fire broke out in building SBG2. Uncontained fire spread and shared power pathways destroyed building SBG2, damaged adjacent SBG1, and forced emergency shutdown of SBG3 and SBG4.

Impact: Thousands of servers were physically destroyed, taking over 3.6 million websites offline worldwide and destroying unbacked-up customer data.

Fix: OVHcloud rebuilt the site with independent physical fire barriers, isolated power distribution per building, and mandatory free off-site backup storage.

Systems lesson: Compute infrastructure extends beyond GPUs and network topology. Fire compartments, power isolation, backup geography, and disaster-recovery assumptions determine whether a rack-scale incident becomes a site-scale outage. Multi-node ML training clusters inherit this physical dependency. A single data-center power or cooling domain failure halts active training runs and destroys un-replicated checkpoint states unless recovery paths cross physical facility boundaries.

Once rack-level controls are in place, the reliability question moves up to facility redundancy, where the same failure-domain logic that isolates a leaking rack also decides how much duplicate power, cooling, and recovery capacity the site should buy. Data center reliability is commonly classified using the Uptime Institute Tier system, with Tier III and Tier IV representing common options for mission-critical compute (Uptime Institute 2026). A Tier III facility is concurrently maintainable: redundant capacity components and distribution paths let operators maintain or replace equipment without shutting down IT operations. Tier IV adds fault tolerance through independent, physically isolated systems, so a single equipment failure or distribution-path interruption should not affect IT operations. The practical difference is not a magic availability number; it is how much failure prevention the building buys in electrical and mechanical infrastructure before software-level resilience takes over.

For ML training, the calculus differs from traditional enterprise computing. A training run is a long-running batch process, not a real-time transaction system. A facility outage during a two-week run therefore usually costs a restart from the last checkpoint rather than permanent data loss, provided a valid checkpoint survives in an accessible failure domain and usable compute becomes available. The restart still incurs lost work, recovery time, and potentially a missed deadline. Many training deployments therefore pair concurrently maintainable (Tier III) facilities with aggressive application-level fault tolerance. Rather than buying every layer of building-level redundancy, they invest in checkpointing, job restart, and storage that can absorb rare facility interruptions, and redirect the avoided resiliency spend toward accelerators, networking, or storage. Facility redundancy and application-level recovery address different parts of the same risk. A design can trade some preventive infrastructure for checkpointing and recovery capacity only after checking which failures the recovery path survives and whether the resulting completion-time risk is acceptable.

Checkpoint 2.3: Infrastructure physics

A team is planning to deploy 256 H100 GPUs (32 nodes, each drawing about 8.4 kW as computed in section 2.6.3) in an existing air-cooled data center with 250 kW of available facility power, a PUE of 1.5, and a cooling plant that can reject 200 kW of IT heat.

- ☐ **Facility power**: calculate an upper bound on the total power draw by applying the PUE to the node figure, and explain why it overstates the draw when the node figure already includes allocated cooling overhead.
- ☐ **Binding capacity**: identify whether power capacity or cooling capacity is the binding constraint.
- ☐ **Required modifications**: specify what modifications would be needed to support the full deployment.

The rack concentrates power and heat into a physical volume where thermodynamics, not software, sets the limits. A single rack of 32 GPUs, however, is far from sufficient for our 175B model, which needs tens to hundreds of nodes, many racks' worth, to train on a practical schedule. The next level of infrastructure aggregates tens to hundreds of racks into a unified computing system: the pod,

whose physical interconnect fabric must be engineered around the communication patterns and failure domains the workload creates.

2.7 The Pod

Training our 175B model on a single DGX H100 node (8 GPUs, 7,912 TFLOP/s of aggregate dense FP16 peak) would take about 2.8 years at the efficiency assumed below, even if the model state could fit. Roughly 73 nodes bring the compute-only time down to two weeks, and about 1,024 nodes bring it to a single day, before any communication or recovery overhead. A pod is the cluster-scale unit that makes this possible: many racks under one high-speed fabric, power envelope, and cooling design. The term names that organization rather than a fixed size, so the pods in this chapter range from the 1,024-GPU reference cluster to the larger configurations worked below. The engineering challenge is wiring those nodes together into a network fast enough to keep gradient synchronization from becoming the bottleneck, so that the fabric serves the role a system bus serves within a single machine. Pod design is the first point where topology, power, and failure domains become shared infrastructure rather than per-rack concerns.

Two fabric families dominate. A switched fat tree, a hierarchy of switch tiers with as much uplink as downlink capacity, supplies full bisection bandwidth across node uplinks, while a direct-connect torus wires each node to its neighbors, using fewer switches but giving bisection bandwidth that grows more slowly than node count. Section 3.5.2 and section 3.5.4 develop both.

The pod-scale examples that follow deliberately step up to eight times the 1,024-GPU reference cluster. A 1,024-node DGX H100 cluster contains 8,192 GPUs, 4,096 NVSwitch chips, 8,192 InfiniBand HCAs, several hundred InfiniBand switches, tens of thousands of cables, and draws approximately 8.6 MW of rack power, including each rack's allocated cooling overhead. It occupies 256 four-node racks across one or more data center halls.

The physical weight of such a cluster is also substantial. Each DGX H100 node weighs approximately 130 kg (288 lb), so 1,024 nodes weigh over 133 metric tons. Racks, switches, cables, and cooling infrastructure add substantially to that total. The data center floor must be structurally engineered to support this concentrated load, which is far denser than the loading that traditional server rooms were designed for. This structural requirement is often overlooked during facility planning and can be a blocking constraint for retrofitting existing buildings.

The physical layout of the data center hall reflects these density constraints. The high power density of ML training racks requires strict hot aisle/cold aisle containment in air-cooled sections, where cold air is forced into the enclosed front of the rack and waste heat is captured immediately at the rear exhaust. The "spine" of the hall, the central cable corridor connecting all racks to the aggregation switches, must accommodate thousands of fiber optic cables and power feeds. Overhead cable trays are preferred over under-floor routing to improve airflow and accessibility, carrying the heavy copper power feeds and fragile fiber interconnects. The layout must also accommodate the liquid cooling infrastructure: coolant distribution unit placement, coolant piping runs, and isolation valves that allow individual racks to be serviced without draining the entire cooling loop. These physical layout decisions, made during facility design, constrain the network topology options available to the training team years later.

Training a 175B model on this cluster costs about $6 \times P \times D$ floating-point operations for P parameters and D training tokens, the training-compute estimate of section 1.3.2. The token count is where the scaling laws of section 1.3 enter the hardware plan, because a compute-optimal allocation would give a model of this size many times more tokens than GPT-3 saw, multiplying this budget, and the pod-hours it consumes, by the same factor. Notebook 2.4 turns the budget into a compute-only baseline under an assumed local compute efficiency. That efficiency is not measured distributed MFU, because MFU over a distributed step already includes its communication and pipeline losses. On the 1,024-GPU reference cluster itself the same baseline would take about 8 days. The Fleet Law (equation 1.3) writes a distributed step as local compute divided among the devices, plus communication and synchronization time, minus the part of that overhead that overlaps computation. The compute-only baseline is the T_{compute}/N term alone.

In practice, communication overhead, pipeline bubbles, checkpoint I/O, hardware failures, and maintenance windows compound to extend the actual wall-clock time beyond that baseline. The first two are the Fleet Law's exposed $T_{\text{comm}}(N)$ and $T_{\text{sync}}(N)$ terms, and the rest are operational

losses around the step. Scaling efficiency, how close a cluster comes to a speedup in proportion to its device count, charges those two terms as a single factor. That factor typically falls as a cluster grows, because exposed communication and synchronization take a larger share of each step, as section 6.4.1 derives. The notebook chains it with the operational losses into a scenario estimate of wall-clock time; production schedules can stretch further once queueing, data readiness, debugging, reruns, and site-specific operational margins enter the plan. At pod scale, the schedule therefore depends on fleet overhead as well as on local execution, and recovering the lost time draws on better networks, overlapped communication, and proactive maintenance.

Napkin Math 2.4: Training time for 175B

The training time for our 175B model derives from first principles.

1. **Total FLOPs**: Using the same $6 \times P \times D$ approximation, where $P = 175 \times 10^9$ and $D = 300 \times 10^9$ tokens: 3.15×10^{23} FLOPs
2. **Cluster throughput**: 8,192 H100 GPUs at the 989 TFLOP/s dense FP16 peak of 16-bit mixed-precision training, assuming 45 percent local compute efficiency before distributed overhead: $8,192 \times 989 \times 10^{12} \times 0.45$ FLOP/s before distributed overhead $\approx$ 3.65 EFLOP/s
3. **Compute-only training baseline**: $(3.15 \times 10^{23}$ FLOPs$)/(8,192 \times 989 \times 10^{12} \times 0.45$ FLOP/s$) \div 3600 \approx$ 24.0 hours
4. **Illustrative overhead chain**: The scaling factor already includes pipeline fill and drain, so no separate bubble step follows. Each subsequent percentage adds time relative to the preceding subtotal, not to final wall time; the factors are scenario assumptions, not independent measured losses:
   - Distributed-step losses (scaling efficiency 0.56, the modeled tier for 8,192 GPUs, covering communication and pipeline bubbles): $\div$ 0.56 $\rightarrow$ 42.9 hours
   - Checkpoint overhead ($\approx 3\%$): $\times 1.03 \rightarrow$ 44.1 hours
   - Hardware failures and restarts ($\approx 10\%$ of the checkpoint-adjusted subtotal): $\times 1.10 \rightarrow$ 48.6 hours
   - Maintenance windows ($\approx 5\%$): $\times 1.05 \rightarrow$ 51.0 hours

The total training time is 51.0 hours in this simplified scenario chain, or about 2.1 days. Real production calendars may still stretch to weeks after queueing, debugging, reruns, and site-specific operational margins are included. Recovering this lost time is the domain of distributed systems engineering.

2.7.1 The pod as a warehouse-scale computer

When a single training job spans 10,000 H100 GPUs across an entire data center hall, about ten times the reference cluster, a slow network tier or a single device failure stalls the whole building's computation. At this scale, the pod operates as a warehouse-scale computer (WSC) (Barroso et al. 2019) and requires a matched budget for arithmetic, communication, and recovery. Rounding the pod up to 10,000 GPUs keeps the warehouse-scale arithmetic easy to follow, and the chapter's other fleet-scale power, cost, and failure examples use the same round figure.

Such a WSC supplies approximately 9.89 EFLOP/s of FP16 peak and 19.8 EFLOP/s at FP8. Non-blocking InfiniBand at 50 GB/s per GPU, the injection bandwidth each endpoint can push into the fabric, provides 500 TB/s of aggregate injection bandwidth and 250 TB/s across the bisection, the cut that splits the cluster into two equal halves. With 4:1 oversubscription at that boundary, meaning the switch tier above it has one quarter as much uplink as downlink capacity, the modeled crossing capacity falls to 62.5 TB/s. The effect on a collective depends on how much traffic crosses the constrained paths, its schedule, and competing transfers. The bandwidth ratio alone is not a collective-time multiplier.

↩ The warehouse-scale computer as the unit of fleet design is established in section 1.2.2.

Treating the building as one computer means that a server assigned to a distributed training job contributes to a shared completion time, and the network fabric, distributed storage, and fleet orchestrator make the building's resources available as one training system. Rack placement, the

power delivery chain, and the network topology must therefore be co-designed, because a mistake at any layer propagates through the others.

The first physical consequence is the communication wall at pod scale, where the network can set the scaling ceiling. Within a single node NVLink supplies enough bandwidth for tensor parallelism, but across nodes the fabric must carry gradient tensors, checkpoints, and pipeline stage outputs, so a building-scale machine is only as fast as the wires that bind its cores. The WSC reframing makes the interconnect a first-class architectural component rather than a peripheral.

The second consequence is that *failure is a physical certainty, not an exceptional event*. Under the scenario assumption of independent failures, the mean interval between GPU failures is the per-GPU MTTF divided by the GPU count. At a per-GPU MTTF of approximately 50,000 hours, a cluster of 10,000 GPUs therefore has a mean interval of roughly five hours, and the 25K-GPU fleet of figure 2.1 about two. That shrinking interval is the reliability wall. It is an expectation, not a periodic schedule, and shared power or cooling faults add correlated exposure. The node failure estimate of section 2.5.5 and the facility controls of section 2.6.4 are the same exposure at smaller scope. Redundant infrastructure, such as dual-path power feeds and redundant switch fabrics, can prevent some interruptions; checkpoint storage and recovery capacity limit lost progress when a running job must stop. Both belong in the physical design. Redundancy keeps many component failures from halting the job, and checkpoints bound the work lost when one does.

The third consequence is that the building blocks combine differently depending on the workload they must carry. A pod is not assembled from a single canonical part list; the interconnect, the switching, and the memory placement are each chosen against the communication and capacity pattern the dominant workload imposes. Three pod designs show the range (table 2.15). An NVIDIA SuperPOD wires GPUs through an InfiniBand fat-tree for any-to-any flexibility, a Google TPU pod hard-wires chips into a torus for dataflow throughput, and a recommendation pod built for Archetype B arranges the same class of accelerators around embedding memory rather than raw compute.

Table 2.15: Pod Architecture by Workload: Each pod's interconnect, switching, and memory emphasis reflects the communication and capacity pattern of its dominant workload, not a single best design.

System	NVIDIA SuperPOD	Google TPU Pod	Recommendation pod
Interconnect	InfiniBand Fat-Tree	3D Torus ICI	RoCE (RDMA over Ethernet)
Switching	External Switches	Direct Chip-to-Chip	Leaf/Spine Ethernet
Topology	Any-to-any	Nearest-Neighbor Mesh	Rail-optimized
Optimization	Flexibility	Dataflow Throughput	Embedding Memory
Dominant workload	Diverse Research	LLM Training	Recommendation

The recommendation pod is the clearest case of memory placement, not compute, driving the layout, and it is the physical form of Archetype B (DLRM at Scale). Its embedding tables reach 10 TB or more, which lighthouse 2.1 showed would need at least 15 DGX H100 nodes of HBM merely to hold, while the dense neural-network layers that operate on those embeddings are comparatively small. A recommendation-first design splits the two across the memory hierarchy. The terabyte-scale embedding tables reside in host CPU DRAM (roughly \$3–\$5/GB), and the dense layers execute on GPUs, where the compute density justifies the \$10–\$15/GB cost of HBM. Embedding tables create a capacity requirement, but small random lookups can also saturate aggregate memory bandwidth or expose access latency at high concurrency. Host placement trades capacity cost against lookup throughput and host-to-accelerator transfer, so the workload must validate both budgets. The systems lesson is that optimal infrastructure is workload-specific, not specification-maximizing. Deploying compute-optimized nodes for a memory-capacity-bound workload would strand most of the accelerator's throughput on a problem its silicon was never the bottleneck for.

The physical design supplies constraints, not an operating policy; later chapters turn its power, capacity, topology, reliability, and sustained-throughput measurements into placement, recovery, monitoring, and ownership decisions. One hardware decision remains for this chapter. With every

level from the accelerator to the warehouse-scale computer in view, the choice of accelerator can now be judged against the whole hierarchy rather than against its data sheet.

↪ The ownership decision, in which sustained fleet utilization sets the build-versus-buy threshold, is developed in section 12.9.2.

2.8 Accelerator Selection for Fleet Workloads

The same H100 that sustains 30–50 percent of peak throughput during well-tuned large-batch training may achieve less than 5 percent utilization during single-request inference, because the binding constraint shifts from arithmetic to memory bandwidth. An accelerator's compute throughput, memory bandwidth, memory capacity, and power consumption are meaningful only in relation to the workload running on it. Selecting the right accelerator therefore requires mapping each workload archetype to its position on the roofline plot and identifying which physical resource dominates its performance.

For the 175B running example, large local matrix operations can place substantial training work above the roofline ridge point. Arithmetic throughput then matters alongside the capacity needed for sharded state and the interconnect needed by the partition. H100, B200, and TPU v5p configurations are candidates to evaluate under the same workload and quality target. The selected software stack must express the required operations efficiently enough for the hardware advantage to appear in sustained execution.

Serving the same language model changes the answer, for the reason the token-latency notebook 2.1 established. At batch size 1, the H100's 3.35 TB/s bandwidth determines decode throughput while much of the arithmetic silicon waits, so the key metric becomes memory bandwidth per dollar rather than peak TFLOP/s. Batching helps for the reason section 2.3.2 quantified, reusing each weight read across more requests, which is why serving systems work so hard to batch without violating tail-latency service-level objectives (SLOs). The workload did not change names, but the binding term of the single-machine iron law moved from arithmetic to data movement, and the accelerator choice changes with it.

Recommendation models split the decision across sparse lookup and dense computation, the memory-placement case that the recommendation pod of section 2.7.1 already resolved. Vision workloads and diffusion models, which generate an image by removing noise from it over repeated passes of a denoising network, move back toward arithmetic throughput because convolutions (table 2.7), attention blocks, and denoising networks reuse data heavily. Diffusion inference makes this especially clear. A standard Stable Diffusion v1.5 configuration runs 50 denoising steps, so the service pays for a sequence of full model evaluations rather than one decode step per token. For these workloads, TFLOP/s per dollar often dominates bandwidth per dollar once model state fits in memory.

Mixture of experts (MoE) activates only part of a much larger parameter set for each token (Shazeer et al. 2017); Mixtral-8x7B (A. Q. Jiang et al. 2024) has 46.7B total parameters but only 12.9B active per token because its router chooses 2 experts of 8 experts experts. The router sends tokens to those experts across devices through the AllToAll collective, so the selection consequence is a fabric with enough bisection bandwidth, together with load balancing, for the theoretical saving to appear in the training trace (Fedus et al. 2022; Lepikhin et al. 2021). Multi-modal models similarly combine phases with different arithmetic intensity and batching behavior, such as vision encoding and text decode. Their relative cost and placement determine whether flexible scheduling, additional memory, or communication capacity changes the critical path; modality count alone does not select an accelerator architecture, so the planner must follow the data exchanged between the actual phases.

↪ MoE AllToAll routing and expert load imbalance are developed in section 5.8.1.

> **Systems Perspective 2.2: Matching hardware to workload**
>
> Workload diversity can justify a heterogeneous fleet, with each pool provisioned for its binding resource. The gain comes from matching hardware to the work assigned to it, while keeping each tightly synchronized training group balanced. The fleet-wide decision must account for that matching benefit together with the cost of supporting multiple platforms.

2.8.1 Precision and throughput trade-offs

Booking a twofold training speedup from FP8 on a capacity spreadsheet, or sizing a serving pool for FP16 weights when INT4 would pass quality validation, can misstate fleet device counts by a factor

of two or more. The number formats and peak rates of section 2.4 become fleet decisions as soon as they set how many accelerators a job requires. FP8 doubles the H100's advertised arithmetic ceiling for supported operations, from 989 TFLOP/s to 1,979 TFLOP/s per GPU, but that doubling is a peak-capacity comparison, not a training speedup that can be booked before workload validation. Infrastructure sizing must use the sustained rate of a validated precision policy.

For inference, quantization changes both weight storage and the possible hardware topology. Our 175B model at FP16 requires 350 GB, so the parameters alone span at least five 80 GB H100 GPUs. Quantizing to INT8 halves the weight payload to 175 GB, which fits on three GPUs, and INT4 shrinks it further to 87.5 GB, which fits on two GPUs. These are weight-only capacity bounds; quantization metadata, attention state, and workspace also need memory. Reading 4-bit rather than 16-bit weights quarters the raw payload, but it does not change physical bandwidth or guarantee a fourfold token-latency reduction. Kernel support, conversion work, other traffic, and any change in device count determine the realized gain, subject to validation of model quality.

2.8.2 Benchmarking accelerator performance

A procurement team comparing accelerators needs a published number that predicts the performance the organization will actually achieve. Peak TFLOP/s, memory bandwidth, and TDP describe *potential*, not *achieved* performance, so the decision turns on whose measured results to trust and under what rules they were produced. The industry-standard reference is MLPerf, maintained by MLCommons. MLPerf Training reports the time to train reference models to a target accuracy, and MLPerf Inference reports serving throughput and latency under realistic batching and streaming conditions. Both capture complete system performance, including software stack efficiency, communication overhead, and memory management, rather than raw silicon speed alone.

Closed-division results fix the workload and quality target so implementations can be compared, and open-division results allow broader changes, rules that section 9.3.1 reads as one instance of a measurement contract. Neither division predicts a deployment automatically. A team must compare the submission's workload, software, precision, hardware count, and quality target with the intended system before using its result for procurement.

For our 175B model, training, fine-tuning, and serving need separate measurements. Batch shape, trainable state, sequence length, and partitioning can change both capacity demand and the critical path. Fine-tuning does not occupy a fixed point between training and decode merely because it updates an existing model. The team must evaluate the intended phases on candidate hardware with the actual model and data pipeline, using published MLPerf results as comparative evidence rather than deployment forecasts. MFU reports the share of the chosen peak that useful model computation averages over a measurement window. Time to train a reference model to a target metric also captures the I/O stalls, checkpointing, and restart delays that a steady-state MFU window can omit, although neither establishes the model quality reached.

2.8.3 Inference-specific infrastructure considerations

Training and serving place opposite demands on the same cluster because they optimize over different timescales. A large training run is one long job that occupies its allocation for days or weeks and is judged by sustained *throughput* and time to a target quality. Its steps repeat the same exchanges, so communication demand is easier to characterize than an arrival stream, and an individual operation can take milliseconds to seconds without harm, although failures, I/O, and changes in execution shape still affect the schedule. Serving trained models must instead turn the same hardware into responses for many independent, short-lived requests arriving at irregular intervals, each under a latency target that for interactive use is measured in tens of milliseconds.

Serving infrastructure must meet *latency* targets at an acceptable *cost per query* under a variable request load. Independent requests can be assigned to separate replicas, but a replica that shards a model across devices still communicates within each request. Prompt processing, successive decode steps, and complete responses also have different latency budgets. Capacity planning must therefore specify the request mix, model placement, and service targets before concluding that serving needs less network capacity than training; training and serving part ways on utilization, model placement, network traffic, and the cost metric that judges them (table 2.16).

Table 2.16: Training vs. Serving Infrastructure: Training and serving optimize different hardware behaviors, so capacity planning must compare utilization, model placement, network traffic, and cost in the mode the fleet will actually run.

Dimension	Training infrastructure	Serving infrastructure
GPU utilization	Large local batches push dense layers into the compute-bound regime; well-tuned runs commonly sustain 30–50% MFU.	Batch-1 decode is deeply memory bound and can fall below 5% MFU; batching raises utilization at the cost of latency, within memory capacity.
Model placement	Tensor and pipeline parallelism distribute a single model across many GPUs.	Replicas of the model, or large shards of it, handle independent requests, so memory demand scales with concurrent serving capacity.
Network traffic	The partition determines gradient, activation, and parameter exchanges between workers.	Independent replicas can serve separate requests, but sharded replicas still require inter-device exchanges within each request.
Cost metric	Cost per completed run at a target quality includes the measured schedule and recovery overhead.	Cost per request or per 1,000 tokens depends on request shape, latency targets, utilization, and provisioned capacity.

Serving therefore cannot be planned as training with smaller batches, because the optimization target changes from synchronized throughput to latency-bounded, continuously billed requests. The implication is that effective serving infrastructure often uses different hardware, different software, and different facility designs than the training infrastructure. Some organizations use previous-generation GPUs (A100s) for serving because the memory bandwidth per dollar can be competitive with newer GPUs, and the lower TDP (400 W vs. 700 W) allows higher rack density and lower cooling costs.

The inference workload itself splits into two distinct phases with opposing bottlenecks. The prefill phase processes prompt tokens in parallel, increasing weight reuse and potentially approaching the compute ceiling when the prompt and batch shapes supply enough work. The decode phase generates tokens sequentially for each request and is deeply memory-bandwidth-bound at small batch sizes. Serving systems raise utilization by admitting and retiring requests at iteration boundaries, called continuous batching, and by limiting how much prompt work delays the next decode iteration, mechanisms that chapter 10 develops.

A napkin estimate shows how quickly serving capacity reaches training scale. One tensor-parallel replica of the running model on an eight-GPU node streams its FP16 weights once per decode step, which takes at least 13.1 ms at H100 HBM bandwidth. Batching 32 concurrent requests yields about 2,450 tokens/s per replica. A service handling 1,000 requests/s at 250 output tokens each must generate 250,000 tokens/s, which requires at least 103 replicas, or 824 GPUs, most of the 1,024-GPU training reference cluster. The estimate is a lower bound because it ignores key-value cache traffic, prefill work, and headroom for bursts and failures; section 10.6 replaces it with batching models that account for key-value cache capacity, variable request lengths, and latency targets. Reliability objectives differ as well. A training job can absorb an occasional failure because checkpointing recovers the lost progress, whereas a stalled serving replica delays every request routed to it until traffic fails over.

2.8.4 The accelerator decision matrix

Combining workload regime, precision policy, and serving latency targets yields the procurement screen in table 2.17. Its candidate configurations rest on representative assumptions rather than fixed assignments of model families to bottlenecks or vendors. The rule is to diagnose the binding constraint first, then compare sustained useful work within the capacity, quality, latency, and facility budgets. The table's constraint column is the per-device one, what a single accelerator must supply. The archetype roster of section 1.6.1 answers a different question, which constraint binds the fleet as a whole, so a workload can be compute-bound on each device and still communication-bound across the cluster, as Archetype A is.

Table 2.17: The Accelerator Decision Matrix: The optimal accelerator depends not on peak specifications but on which physical resource (compute, bandwidth, or capacity) is the binding constraint for the target workload. The Roofline Model screens the compute and bandwidth ceilings at the intended precision and shape. Among candidates that clear capacity and facility constraints, measured sustained useful work at the required quality and latency determines value.

Workload	Binding constraint	Key metric	Recommended class
LLM training (>100B)	Compute (high batch size)	Sustained useful throughput at validated precision, interconnect BW	H100/B200, TPU v5p
LLM inference (batch=1)	Memory bandwidth	GB/s per dollar	A100, H100 (bandwidth/$)
LLM inference (batched)	Compute + bandwidth	Sustained token throughput within latency and capacity limits	H100, B200
Vision model training	Compute (large spatial dims)	Sustained useful throughput at validated precision	H100/B200 configurations to evaluate
Recommendation (embeddings)	Memory capacity	GB per dollar	CPU DRAM + GPU hybrid
Fine-tuning (<13B)	Memory capacity (state near HBM limit)	HBM capacity	A100-class capacity candidate
Research/prototyping	Flexibility	Software ecosystem	GPU or accelerator with required operator support

For our 175B running example, table 2.17 separates the training purchase from the serving purchase. The workload constraint comes first; software support, deployment lifetime, and measured sustained throughput are compared only among candidates that clear it. A cheaper accelerator-hour is useful only if it supplies the required throughput within the capacity, latency, and facility budgets.

Checkpoint 2.4: Accelerator selection

An engineering team needs to deploy a 70B-parameter model for both training and inference. Training will use batch size 2048 across 256 GPUs for 3 months. Inference will serve 10,000 requests per second at batch size 1 for 2 years.

- ☐ **Training intensity**: identify the local microbatch, sequence length, and partitioning information needed to estimate training intensity; compare it with batch-1 decode under the chapter's weight-traffic approximation.
- ☐ **Hardware split**: identify which measurements are needed to determine whether training and inference should use the same hardware or different hardware.
- ☐ **Two-year cost**: formulate the 2-year inference-cost comparison between H100s rented at a cloud rate of \$4/GPU-hour and A100s already in the fleet, charged at its reference in-service rate of \$2/GPU-hour, and identify the prompt and output lengths, replica placement, sustained throughput, and capacity assumptions required to evaluate it.

Each of these selections assumes the memory, communication, and power boundaries stay fixed at today's package and rack limits. Hardware technologies now entering deployment do not remove those walls, but they relocate where each constraint binds across a facility's lifetime.

2.9 Infrastructure Technologies That Move the Walls

A new data center will host several generations of accelerators over its facility life, which operators typically depreciate across 15–20 years (Barroso et al. 2019), so future-facing technologies matter only when they change a durable planning constraint. The useful test is not whether a technology is novel, but which wall it moves, what new wall appears after the move, and which fleet-design decision changes as a result. Compute Express Link (CXL) moves the memory-capacity boundary, wafer-scale integration the chip-to-chip communication boundary, and optical circuit switching the fabric topology and power boundary, while advanced packaging moves the die boundary inside the accelerator package. None of them repeals the physics of data movement; each relocates the place where infrastructure teams must spend power, cooling, cost, and software complexity.

2.9.1 Compute express link (CXL)

Holding 1,400 GB of Adam optimizer state or terabyte-scale embedding tables inside accelerator HBM consumes the most expensive memory tier in the cluster for data touched only once per parameter update or sparse lookup. **CXL** addresses this capacity bottleneck by building protocols for I/O, caching, and coherent memory access on top of the PCIe physical layer. Supported device types and platform configurations determine which of those capabilities are available. Memory expanders and pooled memory can make host capacity less rigid, but accelerator access also requires a supported data path and software placement. Fleet planning must establish those capabilities before treating a pool as available training memory.

For ML training, the natural target is cold or phase-local state rather than the active matrix-multiply working set. The 1,400 GB Adam optimizer state of our 175B model is touched during parameter updates rather than streamed through every Tensor Core operation. A shared CXL pool could therefore reduce per-node memory pressure by holding optimizer shards, embedding-table spillover, or local data-cache buffers outside the accelerator HBM tier. The training framework still has to schedule this state carefully, but the placement problem becomes less binary than "fits in this node" or "spills to storage."

The new wall is bandwidth. A representative CXL memory path over PCIe Gen5 x16 provides 64 GB/s, while H100 HBM provides 3.35 TB/s. That 52× gap means weights and activations for the forward and backward passes still belong in HBM. CXL is a capacity tier, not an HBM replacement. For fleet design, planners should ask whether server platforms, rack layouts, and orchestration software can absorb CXL memory pooling, not whether CXL removes the need for high-bandwidth accelerator memory.

2.9.2 Wafer-scale planning constraint

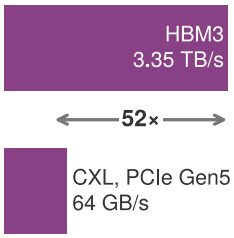

HBM remains 52× faster, so CXL is capacity only.

At the opposite extreme from CXL sits the wafer-scale engine introduced in section 2.2.1. Its on-wafer fabric removes package- and board-level transfers for mapped operations, but it does not make communication free. Infrastructure planning instead shifts toward capacity, placement, and injection bandwidth.

Our 175B running model needs 350 GB for FP16 weights alone, far beyond the 44 GB of on-wafer SRAM. Cerebras addresses this through weight streaming: external MemoryX units feed weights to the wafer over a 1.2 TB/s path while the wafer executes each layer, decoupling model capacity from the on-wafer memory. Weight streaming was designed for training. Activations stay on the wafer while weights stream in layer by layer and gradients stream back to MemoryX, which performs the optimizer update. The approach works when each layer's computation is large enough to hide the streaming time. The wafer trades the communication wall for a streaming-bandwidth wall.

The lesson for infrastructure planners is not that wafer-scale systems replace GPU fleets. It is that memory proximity and memory capacity trade against each other. WSE-style designs buy extraordinary local bandwidth by putting the working set near the arithmetic, while CXL buys flexible capacity by moving memory farther away. HBM sits between those extremes: off-chip but on-package, smaller than host memory but fast enough for the active computation. The right architecture is the one whose memory tier matches the workload's reuse pattern.

2.9.3 Reconfigurable topologies and optical circuit switching

Beyond the package, optical circuit switching (OCS) moves the fabric topology and power boundary. Micro-electro-mechanical systems (MEMS) mirror arrays steer light between fiber ports without converting it to electrical form. A pod's physical topology can therefore be rewired for each job and around failed hardware, as in Google's TPU v4 pods. Millisecond-scale reconfiguration, however, limits OCS to job-level allocation rather than per-message routing. Section 3.5.5 develops the mechanism, the TPU v4 deployment, and its fault-isolation limits.

The common theme is relocation rather than elimination. CXL pushes capacity outward into a fabric where PCIe bandwidth becomes the limiting discipline, whereas wafer-scale integration pulls communication inward onto silicon where on-wafer capacity and yield take over. In the same way, advanced packaging, with 2.5D interposers such as TSMC's CoWoS and 3D stacking such as SoIC, compresses die-to-die distance past the reticle limit. It leaves the rack to absorb the resulting power density, which must be checked against the package and cooling budgets. Optical

circuit switching, in turn, removes the optical-electrical-optical conversion inside the switch only to make reconfiguration latency the new boundary. Infrastructure planning therefore still begins by identifying where the next byte moves, how quickly it must arrive, and what physical system must be built so that movement does not dominate the workload.

2.10 Fallacies and Pitfalls

Because the bottlenecks move rather than disappear, compute infrastructure teams often optimize the visible specification while missing the binding physical constraint. Each fallacy and pitfall below turns a node, rack, pod, or facility lesson into an operational warning.

Fallacy: *More GPUs always means faster training.*

Section 1.8 warned against reading a compute budget as a wall-clock prediction. The hardware form of the same error assumes that training time scales linearly with GPU count, so that doubling the GPUs halves the training time. In practice, communication overhead grows with cluster size, and this chapter's modeled scaling efficiencies show how fast.

Amdahl's law bounds speedup when a fixed workload has a fixed serial fraction. Distributed training also incurs communication, pipeline, and I/O costs that can change with partitioning and scale; these are not all an invariant serial fraction. For a 175B-parameter model with 350 GB of gradients, the AllReduce at each training step must combine contributions across the participating group. It can do so through ring or tree exchanges; direct transfers between every pair are not required.

Even for a carefully engineered fabric, the modeled scaling efficiency of the running model falls from 0.87 at 256 GPUs to 0.83 at 1,024 GPUs and 0.56 at 8,192. Growing from the reference cluster to the 8,192-GPU scenario of section 2.7 therefore multiplies the hardware 8× but the speed only 5.4×, and 44 percent of the larger cluster's locally achievable throughput is lost to communication and pipeline bubbles. That run still finishes in days, so the scale is worth buying for a deadline, but each added GPU returns less than the one before it; without an engineered fabric and overlapped communication, the curve flattens far sooner.

The correct approach is to profile useful compute time, exposed communication time, and idle time at the current scale before buying the next pod. If the added GPUs mostly increase exposed communication, procurement is buying idle time rather than training progress. Section 6.4.1 turns this local ratio into a scaling-efficiency estimate before committing to procurement.

Pitfall: *Planning training throughput from peak FLOP/s alone.*

Peak FLOP/s describes only one roofline ceiling. At the H100 FP16 ridge point of approximately 295.2 FLOP/byte, the bandwidth and arithmetic ceilings meet; lower-intensity operations cannot reach the arithmetic peak through more compute alone. In the chapter's batch-1 weight-streaming estimate, decode near 1 FLOP/byte uses well under 1 percent of the arithmetic peak. Prefill and training supply more reuse, yet even well-tuned large training runs commonly sustain only 30–50 percent of peak once memory-bound attention and normalization kernels, activation recomputation, dependencies, and communication are counted. Procurement needs those measurements for the intended execution phase.

Fallacy: *Power and cooling constraints can be handled after accelerator purchase.*

Teams that plan GPU purchases from compute requirements without verifying that the target facility can deliver the required power and cooling run into hard physical limits with long lead times. A single rack of four DGX H100 systems requires 33.5 kW of power. Upgrading a data center's electrical infrastructure (new transformers, switchgear, and UPS capacity) and its cooling plant (chillers, piping, pumps) are construction projects planned in months, not days. Even seemingly simple changes, such as upgrading the PDU in a single rack from 20 kW to 60 kW, can require new cabling and circuit breaker panels. Organizations that order hardware without confirming facility readiness risk having millions of dollars of GPUs sitting in shipping containers in the parking lot while the building is retrofitted. Accelerators can arrive faster than a facility can be upgraded to power and cool them, which makes this mismatch a practical planning risk.

Pitfall: *Buying a newer accelerator without checking the workload's arithmetic intensity.*

Suppose a team deploys B200 GPUs (4,500 TFLOP/s FP8 peak) for an inference workload serving a 7B-parameter model at batch size 1, where an arithmetic intensity of approximately 1 FLOP/byte places the workload deep in the memory-bound region. The B200's memory bandwidth (8 TB/s) is 2.4× higher than the H100's (3.35 TB/s). Comparing B200 FP8 against H100 FP16/BF16-class throughput gives a peak arithmetic ratio of about 4.6×, but also changes the precision convention. At a fixed weight footprint, the bandwidth-only model instead predicts a gain of about 2.4×. The low intensity limits the fraction of arithmetic peak that this model can use, though it does not measure idle time; the comparison that matters is bandwidth per dollar at the measured serving rate, and the larger arithmetic peak alone does not establish better value.

Fallacy: *Air cooling is sufficient for AI workloads.*

Operators of legacy facilities may assume that adding fans and chiller capacity is sufficient for any denser AI rack. In practice, the cooling choice depends on the system's heat-rejection design and the facility's available airflow, pressure, inlet temperature, and containment. Engineered air remains viable for some dense configurations, while higher rack densities can require direct liquid cooling to avoid thermal throttling or exceeding facility limits.

Fallacy: *All GPU-hours are equivalent.*

A common error in cost comparisons is to treat one GPU-hour on an H100 as interchangeable with one GPU-hour on an A100 or V100. Under the same dense FP16 peak convention, H100 offers 3.2× the arithmetic rate of A100 and 7.9× that of V100. If a workload sustained the same fraction of those peaks with the same operation count and target quality, 10,000 V100 GPU-hours would correspond to approximately 1,264 H100 GPU-hours. That is a conditional compute-only comparison, not a predicted training time, because memory, communication, numerical recipes, and software can change sustained efficiency. Procurement must compare measured cost per useful computation or per completed run, rather than rental price or peak FLOP/s alone.

Pitfall: *Upgrading the network before identifying the training bottleneck.*

A network upgrade from 200 Gb/s to 400 Gb/s InfiniBand changes training time only through the communication it accelerates on the critical path, the mirror image of the faster-accelerator fallacy in section 1.8. In a no-overlap example with 20 seconds of computation and 2 seconds of communication, doubling transfer bandwidth reduces the estimated step from 22 to 21 seconds, an improvement of 4.5 percent. That result assumes the communication term scales with bandwidth and other terms remain fixed. Adding GPUs or optimizing kernels requires a separate estimate because either can change communication and overlap. The Fleet Law keeps those costs visible when comparing investments.

Fallacy: *Average power draw is sufficient for cooling design.*

Accelerator power draw varies across communication, memory, and compute phases. A workload average below the configured sustained envelope does not establish the cooling requirement because synchronized compute phases can raise many devices toward that envelope together. Thermal design must use the vendor's sustained design target under the intended inlet and cooling conditions, while the electrical path must separately tolerate documented transients. Otherwise thermal throttling can reduce sustained throughput without producing a crash or explicit application error.

Pitfall: *Selecting accelerators by HBM capacity without checking bandwidth.*

Capacity and bandwidth answer separate feasibility questions, and a model that fits in HBM but cannot be served within its latency target is as unusable as one that does not fit. Consider two hypothetical options for a 70-billion-parameter model: (a) 192 GB of HBM2e at 2.04 TB/s, and (b) 80 GB of HBM3 at 3.35 TB/s. Option (a) holds the 141 GB FP16 weight payload, whose streaming floor is 141 GB/2.04 TB/s = 69.2 ms. Option (b) cannot hold 141 GB in its 80 GB capacity, but the 71 GB INT8 payload fits before other live state is counted. Its weight-transfer floor is 71 GB/3.35 TB/s = 21 ms, a 3.3× ratio between the two estimates. That ratio combines faster memory with fewer bytes; option (a) could also use a validated INT8 recipe. Selection must compare like-for-like quality, complete live-state capacity, and measured latency at the intended concurrency.

Fallacy: *Data loading can be fixed after deployment.*

Teams that optimize only GPU kernels and communication can still find their accelerators waiting for input. The data loading pipeline must sustain the cluster's consumption rate across storage, preprocessing, and host-to-GPU transfer. The chapter's pre-tokenized language-model example needs modest raw bandwidth, but tokenization, sequence packing, and shuffling can change the host-side demand. Teams must size CPU workers, buffering, and storage concurrency from the actual pipeline. Local NVMe staging and prefetching are useful when they relieve the measured bottleneck; neither substitutes for testing the complete path at fleet concurrency.

Pitfall: *Buying homogeneous clusters without modeling workload diversity.*

Uniform hardware simplifies support and placement, but that benefit must be weighed against the useful work existing devices can still perform. For our 175B model, training and low-batch decode stress different resources. Separate pools can exploit that difference without placing mismatched devices in the same tightly synchronized group.

The A100's 2.04 TB/s and the H100's 3.35 TB/s supply different bandwidth ceilings. Retaining older devices for an inference pool can reduce cost when their measured throughput, power, and support costs compare favorably at the required latency. The distinction is between *job-level* balance within a synchronized group and *cluster-level* hardware uniformity. A heterogeneous fleet can preserve balanced groups, but its matching benefit must cover the added scheduling and operational complexity.

Pitfall: *Treating commissioning and burn-in as schedule slack.*

Commissioning is part of the system build, not a buffer that can be consumed when hardware delivery slips. Burn-in tests, cable validation, firmware qualification, thermal soak, and collective-communication stress runs are the first time the node, rack, network, and facility layers operate as one machine. Compressing that phase turns installation defects into production failures, where each debugging cycle idles the full fleet rather than a small acceptance-test window.

2.11 Summary

Training our 175B-parameter model exposes a constraint cascade: accelerator capacity and throughput shape the node, node power shapes rack cooling, rack density shapes pod layout, and pod topology and reliability determine how much useful work the fleet sustains. At the device level, the generality tax (principle 1) explains the trade between flexibility and workload-specific efficiency. A GPU delivers 10–100× the matrix throughput of a CPU on suitable workloads, but that specialization pays only when the workload and software stack can use the dedicated arithmetic over the deployment's lifetime. The roofline in section 2.3.2 identifies the binding compute or memory ceiling, while the bandwidth hierarchy in section 2.5.1 guides where communicating work belongs. Together, these models connect software choices to the physical limits of the fleet.

Our 175B-parameter model connects the decisions at each level. Its low-batch decode estimate requires about 350 billion floating-point operations per token, but moving the 350 GB FP16 weight payload through one hypothetical 3.35 TB/s HBM path takes over 100 ms. A feasible partition changes that path budget as well as solving capacity. At the node level, training state exceeds 2 TB, requiring partitioning across nodes or reductions in resident state. At the rack level, 32 accelerators at 700 W each, plus modeled host and network overhead, require 33.5 kW of qualified power and cooling capacity. The reference cluster coordinates 1,024 GPUs, while the worked pod scenario uses eight times that number. Communication, failure recovery, and sustained utilization determine how much of the purchased capacity becomes completed work. The model's size exposes these interactions, but the same accounting applies whenever a workload crosses physical resource boundaries.

The method transfers beyond this model. At each level, an engineer identifies the binding constraint, quantifies its effect, and propagates the consequence upward. A faster accelerator is counterproductive if cooling cannot remove its heat; a denser rack is futile if the grid cannot supply it; and a wider network is wasteful if the workload cannot use its bandwidth. Infrastructure engineering is thus a quantitative discipline for tracing chip-level choices into rack power, facility cooling, network topology, reliability, and ownership cost.

Key Takeaways: The data center is the computer

- **Infrastructure walls cascade**: Memory, communication, power, and reliability walls appear at escalating physical scales, each level's limit becomes an input to the next, and new technologies relocate a wall rather than remove it. With independent failures and a per-GPU MTTF of 50,000 hours, the mean interval between GPU failures shrinks with fleet size, to about five hours in the 10,000-GPU example, so recovery and checkpoint storage belong in infrastructure planning from the start.
- **Capacity and bandwidth are separate tests**: Capacity decides whether model state fits; bandwidth decides how fast low-reuse phases run. For single-request LLM decode, weight traffic commonly sets the latency floor. The ridge point (I_{ridge} = peak FLOP/s divided by memory BW) separates bandwidth- from compute-bound operation, but batch shape, fusion, precision, reuse, and the modeled memory level determine where each training or inference phase lies.
- **Bandwidth hierarchy guides parallelism placement**: The 9× per-direction bandwidth gap in the NVLink–InfiniBand example favors keeping frequent tensor-parallel exchanges inside the fast interconnect domain. Pipeline and data parallelism can span slower links, but their traffic and synchronization still have to fit the workload's time budget.
- **A node aggregates private memories**: Eight accelerators multiply HBM capacity, but each device's memory stays private, so the running model's 2.8 TB of persistent training state must be partitioned across nodes. Tensor parallelism inside the node and eight pipeline stages across nodes bring it to 43.8 GB per GPU without sharding, and optimizer sharding, activation checkpointing, and memory tiering are the next steps when activations and workspace do not fit.
- **Plan on sustained throughput, not peak**: A run planned from peak FLOP/s needs GPU-hours in inverse proportion to its measured utilization, so capacity planning must use sustained throughput measured with a consistent model-FLOP convention, precision, system scale, and workload shape.
- **Rack power and cooling bound installable compute**: Watts per rack, delivery losses, and cooling class, not FLOP/s, set how much compute a facility can install and sustain; at pod scale, network topology, cooling, and power delivery must be co-designed for the dominant workload.
- **Failure exposure grows with fleet size**: With independent failures and a per-GPU MTTF of 50,000 hours, the 10,000-GPU example has a mean interval of about five hours between GPU failures. Recovery and checkpoint storage therefore belong in infrastructure planning from the start.
- **Price useful work, not hardware**: A sound purchase diagnoses the binding constraint first, then compares GPU-hours from different generations, an FP8 peak that doubles advertised capacity, and bandwidth per dollar for decode only through sustained throughput on the target workload. Power, cooling, facilities, and utilization then set the lifetime cost of that useful work, so hardware price alone cannot settle a purchase.

Applied through C³, choosing liquid cooling for a dense accelerator rack changes more than heat removal. It raises the compute a rack can hold, which concentrates communication onto fewer top-of-rack links and adds coordination work for power capping, failure isolation, and maintenance. In the Fleet Law's accounting, the hardware in this chapter fixes the per-device rate behind T_{compute}/N and the link bandwidths behind $T_{\text{comm}}(N)$, while rack and pod reliability decide how often synchronization and recovery interrupt the step. Removing one physical ceiling therefore requires checking the rest of the fleet again; otherwise the investment creates capacity that the workload cannot use.

What's Next: From silicon to wires

Physical infrastructure spans from Tensor Core transistors to megawatt data-center substations. A fleet of disconnected nodes, however, is not yet a computer; it is a warehouse of expensive accelerators. The network fabric turns those nodes into a unified training system by carrying gradients, activations, and checkpoints between them. At every scale above the node, this chapter met the same limit in a new form. The bandwidth hierarchy showed a 9× cliff between NVLink and InfiniBand, and the pod showed why usable bandwidth depends on the communication pattern and where it is placed. Every training step of our 175B model ends by combining gradients across all 1,024 GPUs of the reference cluster, so the slowest required transfer sets when the next step can begin. Chapter 3 starts from that synchronization and climbs from signaling on a single link through transport protocols, switch topology, and congestion behavior to the integrated cluster. At each level it asks whether the fabric lets those GPUs train as one coherent machine or leaves them waiting as isolated devices.

3

Network Fabrics

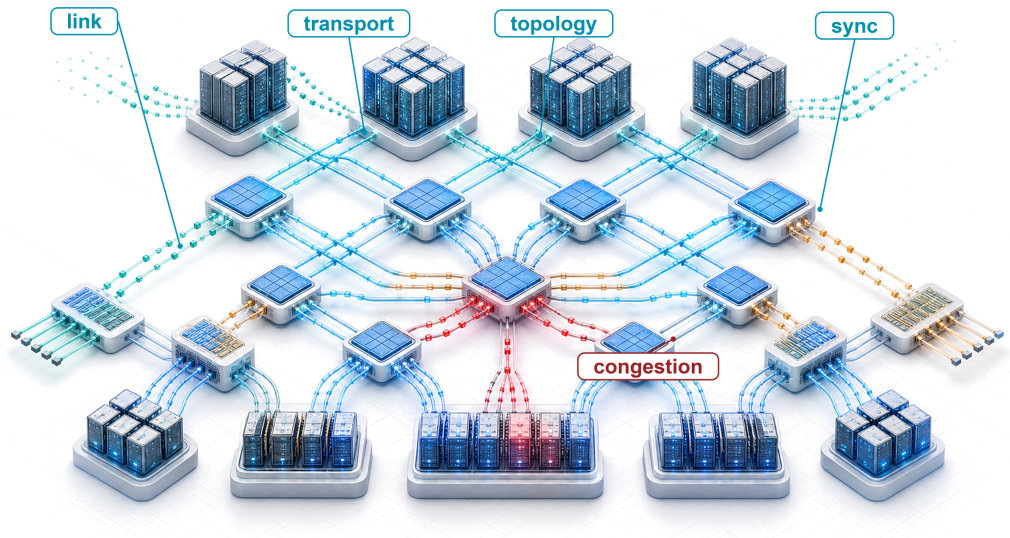

3.1 The Synchronization Backbone
3.2 How ML Networking Inverts Data-Center Assumptions
3.3 Level 1 Wire and Link
3.4 Level 2 Transport and the Performance Model
3.5 Level 3 Switch and Topology
3.6 Level 4 Fabric Behavior, Congestion, and Routing
3.7 Level 5 Cluster Design and Case Studies
3.8 Network Virtualization
3.9 Monitoring and Debugging
3.10 Network Technologies That Move the Ceiling
3.11 Fallacies and Pitfalls
3.12 Summary

Purpose

Why can the network connecting accelerators matter more than the accelerators themselves at scale?

A single accelerator can perform trillions of operations per second, but distributed training requires devices to exchange results before dependent work can proceed. When network capacity cannot keep pace with computation, accelerators sit idle, and adding more devices can increase the waiting instead of shortening the run. The resulting cost extends beyond the network equipment: an inexpensive connection can leave expensive compute capacity unused throughout a long experiment. Single-node intuition is insufficient because data that once traveled through a local interconnect must now cross cables, switches, and shared paths. A fast connection between two machines does not establish that the same performance will hold when the entire fleet communicates at once. The arrangement of those paths determines which groups can exchange data together and where competing transfers accumulate. Large transfers need sustained bandwidth, while frequent dependent exchanges can remain limited by the time required to start and complete each message. A delayed participant can hold up others even when most links appear healthy. Network design therefore has to follow the workload's communication pattern, and validation must reflect the traffic that will actually run. Choices about placement, sharing, and failure recovery affect whether available bandwidth becomes useful progress or recurring delay. In C^3 terms, the fabric connects compute through communication, while coordination determines when the fleet must wait for that communication to finish.

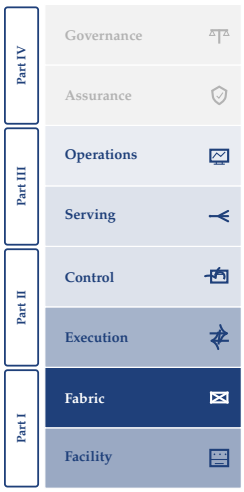

Learning Objectives

- Explain how synchronized collective traffic inverts data-center assumptions about flows, loss, and oversubscription
- Model network communication cost using the α-β framework and identify bandwidth-dominated vs. latency-dominated regimes
- Compare high-performance transport protocols by latency, loss-recovery behavior, and operational complexity
- Compare network topologies using $\text{BW}_{\text{bisect}}$, hop count, and the workload's collective communication pattern
- Evaluate congestion-control mechanisms and lossless-fabric hazards by their impact on tail latency during distributed training
- Compare network partitioning and traffic-isolation mechanisms for checkpoint and multi-tenant traffic
- Diagnose network performance bottlenecks using RDMA counters, link-level telemetry, and bandwidth testing tools

3.1 The Synchronization Backbone

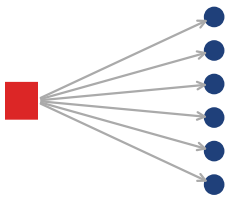

One slow link stalls every peer that waits on it.

At the node boundary, each GPU's NVLink bandwidth narrows to a single external network port, the bottom of the bandwidth hierarchy traced in chapter 2. How much that narrowing costs depends on what the training job must push through the port, and how often. Consider the running instance of Archetype A (GPT-4/Llama-3), a 175B-parameter GPT-3-class dense transformer partitioned across the 1,024 GPUs of the reference cluster. Each training step must combine gradient contributions so that corresponding model replicas apply the same update. This operation, called AllReduce, returns the combined result to every participating worker.

The model has 350 GB of gradient data at two bytes per parameter, the 16-bit gradients of mixed-precision training that this chapter assumes throughout (FP32 gradients would double the figure). That model-wide figure is not what any one GPU sends. When a reduction partitions the gradient across all 1,024 ranks (the numbered workers of the job), each rank owns a 1/1,024 share, a few hundred megabytes, which the transport examples below round to one assumed gradient-shard size; the job's tensor, pipeline, and data-parallel split (chapter 6) sets each GPU's actual payload. Worked examples that switch to another model, precision, or payload state their parameters explicitly.

A delayed transfer on the step's critical path can leave the other 1023 GPUs waiting at a job-wide synchronization boundary. The fabric's contribution to training time therefore depends on which transfers the next computation needs and how much of their cost remains exposed. In that role the fabric is not auxiliary infrastructure; it is the synchronization backbone that decides whether the cluster trains as one machine.

Within the multi-tier fleet hierarchy framed by figure 1.4, the fabric turns isolated accelerator nodes into a coherent training system. Accelerators, power delivery, cooling, racks, and pods define what each node can compute in isolation; the fabric determines whether the nodes can act together before communication cost overwhelms computation. That design space organizes into a **five-level model** of co-dependent tiers, developed in section 3.2.1.

At scale, communication cost can dominate computation cost. The fleet-law decomposition in section 1.5.4 makes this explicit: the nonoverlapped communication term $T_{\text{comm}}(N) - T_{\text{overlap}}$ determines how much of each training step is exposed to the network fabric. The fabric influences the layers above it by determining whether communication can run while independent computation continues, which collective algorithms are viable, and how network partitions interact with node failures. That concurrency, communication-computation overlap, is what the T_{overlap} term credits, because transfer time that proceeds behind useful arithmetic does not lengthen the step.

↪ How layer-by-layer scheduling achieves that overlap, and how much communication it still leaves exposed, is developed in section 5.10.

A collective is a communication operation performed by a group of workers. Its name specifies the result each worker receives, while an algorithm determines the messages that produce that result. The physical network fabric carries three recurring collective communication patterns:

- **AllReduce**: Combines corresponding values from all participants and returns the same result to each; synchronous training uses a sum, followed by normalization, to average gradient contributions across workers.
- **AllGather**: Collects a distinct portion from each participant and returns the concatenated collection to every participant, as required when reconstructing a sharded parameter tensor.
- **AllToAll**: Lets each participant send a distinct portion to every other participant, an exchange central to expert parallelism when workers host different expert subnetworks and must route tokens to their assigned destinations.

Collective communication algorithms orchestrate these patterns, while the fabric layer supplies the physical wires and switches that carry them. The distinction matters because the fabric's physical properties (bandwidth, latency, and topology) determine which patterns are efficient and which become bottlenecks.

↪ The full set of collective primitives, with the output contract each one guarantees, is developed in section 5.4.1.

> Systems Perspective 3.1: The network as a gradient bus
>
> In a single machine, the memory bus moves data between the processor and memory. In a distributed training cluster, the network fabric serves an analogous role as the gradient bus that carries contributions between workers. When those exchanges limit progress, the job encounters the network wall mapped across the pod tier in figure 2.1. Protocols, topologies, congestion control, and collective schedules determine how effectively the job uses the available paths. This bus has a narrow segment. Since NVLink appeared, the external ports that join machines have offered less bandwidth than the interconnect inside each machine in every generation (figure 3.1). How much of that gap a job actually feels depends on its message sizes, port provisioning, routing, and contention, so the link ratio bounds the problem rather than predicting a transfer time.

↩ Communication dominance at scale, the network wall, is established in section 1.2.3.

The reference hardware gives the gradient bus analogy a concrete scale. A single H100 has a precision-specific FP16 peak of 989 TFLOP/s and memory bandwidth of 3.35 TB/s. Each GPU's NVLink ports together provide 900 GB/s of bidirectional bandwidth to its peers within the node. The comparison that matters is per GPU and per direction. In the reference node each GPU has its own network interface card (NIC) with one NDR InfiniBand port, so a tensor leaving the node through that GPU drops from 450 GB/s of NVLink bandwidth to 50 GB/s at the port. This 9× gap is the node boundary's cliff that the bandwidth hierarchy of section 2.5.1 already identified. The model-wide 350 GB gradient state does not all pass through every 50 GB/s port; partitioning and the collective schedule determine each port's traffic. Converting that gap into a step-time estimate requires the workload's local payloads and communication dependencies.

↪ The network hierarchy ratio between intra-node and inter-node bandwidth, with the reference NVLink and InfiniBand values behind it, is developed in section C.2.

Which traffic must cross that gap depends on how the job is partitioned. The same hierarchy keeps tensor parallelism, with its per-layer exchanges, inside the NVLink domain and lets pipeline stages and data-parallel replicas span nodes, the combination called 3D parallelism whose schedules are developed in section 6.6.1. The job has one further lever over that traffic. Each data-parallel replica can sum gradients over several microbatches, the small slices of a batch, before an exchange (gradient accumulation), reducing how often it must cross the node boundary at the cost of changing the batch and update schedule.

Exposed communication leaves gaps between dependent compute phases. Two assumed training schedules on a common 230 ms axis, each containing 100 ms forward and backward compute phases, make that cost visible (figure 3.2). The 8-GPU intra-node case assumes a 1 ms AllReduce and spends 99.5 percent of the step in those compute phases; the 64-GPU inter-node case assumes a 30 ms AllReduce and spends 87 percent in computation. The two AllReduce durations are illustrative inputs chosen to show the effect, not values derived from the link rates above, and the percentages are shares of the schedule rather than measured arithmetic utilization.

In Fleet Law terms, the widened red block is the part of $T_{\text{comm}}(N)$ that no T_{overlap} has hidden. In the inter-node schedule it is 30 ms of every 230 ms step, about 13 percent, and that share applies to every accelerator-hour the cluster spends for the life of the run. Network fabric design therefore belongs at the center of distributed-training design rather than at its edge, and keeping that exposed

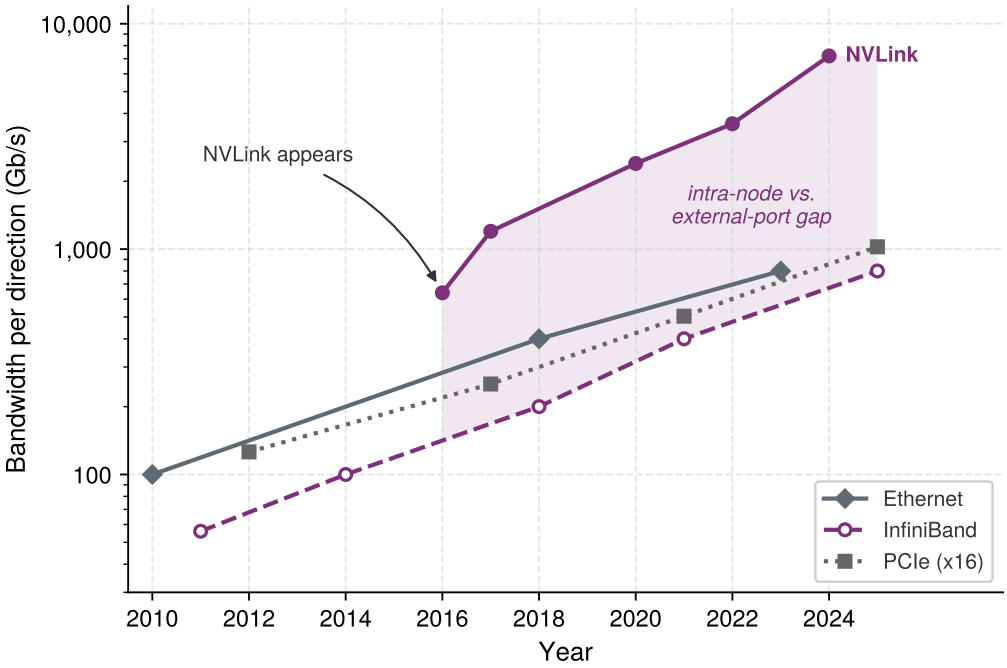

Figure 3.1: Historical Network Bandwidth Scaling: Intra-node NVLink, inter-node InfiniBand and Ethernet, and the PCIe host bus, plotted per direction on a common log scale by the year each generation first shipped, 2010 to 2025. The aggregate NVLink budget exceeds the external-port budgets in these configurations; collective completion time also depends on payloads, provisioning, routing, and overlap. Rates and first-product years from IEEE 802.3, InfiniBand Trade Association, PCI-SIG, and NVIDIA specifications and product announcements.

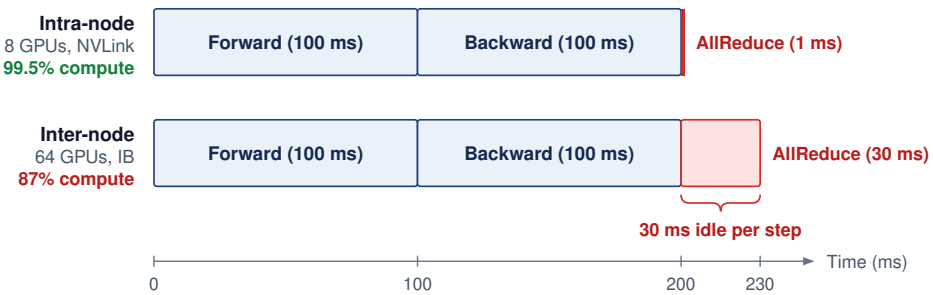

Figure 3.2: The Bandwidth Cliff in a Training Step: Illustrative timelines on a common millisecond axis for the same compute phases: an 8-GPU intra-node job (top) with an assumed 1 ms AllReduce, and a 64-GPU inter-node job (bottom) with an assumed 30 ms AllReduce. These durations illustrate an exposed communication cost rather than predict a particular fabric. The AllReduce phase (red) grows from a thin sliver to a larger share of the step. Without overlap, model computation waits during AllReduce. The compute phases occupy 99.5 percent and 87 percent of these illustrative schedules, a gap of 12.5 percentage points; these shares are not arithmetic utilization measurements.

window small requires rethinking the traffic assumptions that govern conventional data-center fabrics.

3.2 How ML Networking Inverts Data-Center Assumptions

A conventional cloud network built for millions of independent web requests breaks down when a training job releases its gradient exchanges in lockstep at the end of a backward pass. Statistical multiplexing relies on traffic peaks not all coinciding, whereas synchronized workers demand

the same switch paths at the same microsecond. Web services also have dependency chains and tail-latency requirements, but a bulk-synchronous training job ties every participating worker to a single shared completion boundary.

The ML networking inversion in table 3.1 concerns this change in traffic dependencies. Synchronous collective phases can generate correlated bursts and long-lived flows that must complete before the next computation can start. Fabric validation must therefore measure collective completion under the expected load; average per-flow throughput alone does not establish that the job can meet its step-time budget.

Table 3.1: ML Networking Inverts Traditional Data-Center Assumptions: The comparison contrasts independent-request traffic with synchronized collective phases. Correlated demand and shared completion boundaries change the fabric's capacity and latency requirements.

Workload pattern	Traditional data-center assumption	ML reality
Traffic pattern	Independent requests can arrive asynchronously	Collective phases can release correlated bursts
Flow type	Many short transfers can share available paths	Large collective transfers can create persistent "elephant" flows
Performance metric	Per-request completion and throughput	Collective completion and exposed step time
Loss tolerance	Recovery delays affected requests	Recovery can delay dependent workers at a barrier
Congestion	Statistical multiplexing relies on noncoincident peaks	Shared paths and fan-in can concentrate synchronized traffic

The bulk synchronous parallel (BSP)[1] model, the lockstep at step barriers already met at the scale of one node (section 2.5), makes the dependency in table 3.1 explicit. In the synchronous training example, all 1,024 GPUs must complete the required gradient exchange before the next step can proceed. A 100 ms transfer delay adds 100 ms to that step only if it extends the critical path by the full interval; other computation or communication may hide part of it. At a barrier that waits for the last required result, **tail latency**, the slow end of the completion-time distribution (commonly reported as its 99th percentile, p99), decides when the step ends, so the tail cannot be dismissed as an outlier.

[1] **BSP (bulk synchronous parallel)**: Proposed by Valiant (1990) as a bridging model between parallel hardware and software. Synchronous data-parallel training follows the same barrier structure (Goyal et al. 2017; S. Li et al. 2020).

> Definition 3.1: Bulk synchronous parallel (BSP)
>
> ***Bulk synchronous parallel (BSP)*** is a parallel execution model in which every worker completes a local computation phase, exchanges the data required by the next phase, and then waits at a global barrier before any worker begins the next phase, which makes the slowest participant the pacing constraint for the entire cluster.
>
> 1. **Significance**: The slowest required computation or communication determines when a superstep can complete. Faster workers wait at the barrier, so a localized delay can consume idle accelerator time across the entire participating group.
> 2. **Distinction**: BSP organizes execution into synchronized supersteps; it does not itself prescribe a learning algorithm or a particular exchange pattern. Synchronous data-parallel updates match a single-device reference only under compatible global batches, update rules, and numerical assumptions.
> 3. **Common pitfall**: Removing barriers does not by itself establish faster time-to-quality. Asynchronous training trades waiting against stale updates, and the convergence and wall-clock comparison depends on the workload and update schedule.

↩ Synchronous tight coupling, the shared step barrier that lets the slowest worker pace a training fleet, is established in section 1.2.1.

The flow inversion compounds this problem. Traditional networks are designed for fairness among millions of "mice flows,"[2] yet ML training is dominated by a few "elephant flows" corresponding to the gradient AllReduce. The full 175B model has 350 GB of gradient state at the assumed precision, but a sharded job distributes that state across workers. Each collective operates on its participating

[2] **Elephant and mice flows**: The model-wide gradient size is not the size of one network flow. A collective can split each worker's local gradient shard across channels and transport connections; the resulting long-lived flows can still collide under ECMP when their static hashes select the same path.

group and local payload, and its algorithm determines how much each worker sends. Standard routing such as equal-cost multi-path (ECMP), which hashes each flow onto one of several equal-cost paths, is poorly suited to this traffic pattern. A static hash cannot subdivide one elephant flow across all available links, and it can inadvertently map multiple elephant flows to the same link, creating a hot spot while other links sit idle.

The loss inversion completes the picture. TCP was designed to recover across an unreliable network, and an occasional retransmission delays only the request it belongs to. In a synchronous collective, a lost packet delays the transfer that every waiting worker needs, and some NICs that implement remote direct memory access (RDMA), called RDMA NICs or RNICs, recover by resending everything after the missing packet rather than only the lost region (Go-Back-N-style recovery). ML fabrics therefore aim to make loss rare and recovery bounded, either by preventing buffer overflow at the link layer or by building transports that recover quickly from it. Level 2 (section 3.4.3) develops both approaches.

Correlated demand also makes **oversubscription**, provisioning less uplink capacity at a switch tier than the downlink capacity it serves (written as a downlink:uplink ratio such as 2:1), a workload decision. A fabric with spare capacity under independent requests can become saturated when a collective releases traffic across the same cut. Designing for that traffic requires tracing the constraint from the wire through transport, topology, and congestion control to the job's synchronization boundary.

3.2.1 The five-level model

Provisioning fast NICs or buying a non-blocking (fully provisioned) switch tier in isolation routinely fails to speed up training: a 400 Gb/s link still stalls if host TCP copies starve the PCIe bus, an unaligned rack layout exhausts the copper link budget, or static ECMP hashing piles two elephant flows onto one port of an upper-tier spine switch. Tracing why a collective slows down requires following the packet across five coupled layers, from serializer/deserializer (SerDes) signaling and copper or optical reach at Level 1 to the integration of a production cluster at Level 5 (figure 3.3). Each level is developed at the point where its physical or protocol limit dominates.

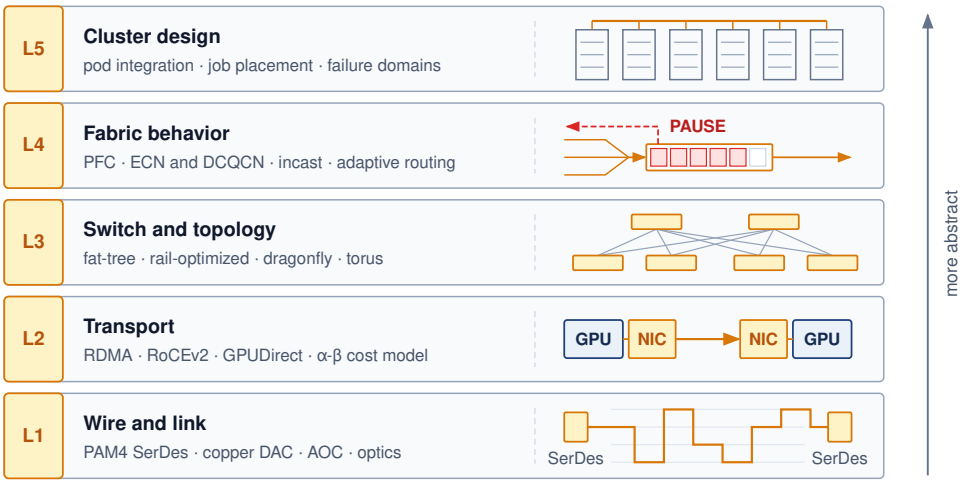

Figure 3.3: The Five-Level Network Model for ML Training: Read from the bottom up, each level builds on the one below and becomes more abstract. Level 1 moves bits over wires and optics; Level 2 turns links into reliable, low-overhead transfers between accelerators; Level 3 arranges switches into a topology; Level 4 governs how the fabric behaves under load and congestion; and Level 5 integrates all of them into a production cluster.

Each inversion of table 3.1 traces back to a specific physical or protocol constraint, and the five levels locate where each one is paid:

- **Level 1: wire and link** (section 3.3). Signal integrity (PAM4, SerDes), forward error correction, and the reach of copper and optics impose hard constraints on per-hop latency and cluster geometry.
- **Level 2: transport** (section 3.4). InfiniBand and RDMA over Converged Ethernet (RoCE) provide RDMA primitives; the α-β model quantifies the latency-vs.-bandwidth trade-off for different message sizes.
- **Level 3: switch and topology** (section 3.5). Fat-trees, rail-optimized designs, and dragonflies achieve the BW_{bisect} (bisection bandwidth, the capacity across the narrowest cut that splits the cluster in half) needed for global collectives through different structural trade-offs.
- **Level 4: fabric behavior** (section 3.6). Link-level flow control and end-to-end congestion control determine whether theoretical bandwidth translates to realized throughput; adaptive routing and incast behavior determine whether that throughput survives real collective traffic (Alizadeh et al. 2010; Zhu et al. 2015; Li et al. 2019; Gangidi et al. 2024).

- **Level 5: cluster design** (section 3.7). Production supercomputers like the NVIDIA SuperPOD and Meta Grand Teton integrate these layers into a unified gradient bus (NVIDIA 2023; Meta Engineering 2024).

The levels are not an inventory of networking topics; they are the causal chain that turns synchronous ML traffic into infrastructure requirements. Large gradient messages first stress the wire and transport, then force topology choices that preserve bisection bandwidth, then expose congestion behavior that ordinary enterprise networks can hide, and finally determine whether the whole cluster behaves like one training machine.

The traffic pattern determines which fabric metric matters. A global reduction stresses bisection bandwidth and the slowest path; a point-to-point pipeline transfer sits on the critical path, so its full transfer time and hop count matter; an expert exchange stresses path diversity and egress balance; and checkpoint or tenant traffic stresses isolation from the synchronization path (table 3.2). The remaining levels develop each of these pairings in turn.

Table 3.2: Traffic Patterns Determine Fabric Metrics: Workload semantics translate into different physical bottlenecks, so no single link-rate benchmark predicts training performance. The table is a map of what follows; the mechanisms in its right two columns (transport offload, adaptive routing, virtual lanes, and the rest) are developed at Levels 2 through 5.

Traffic pattern	Fabric pressure	Metric that predicts performance	Primary design response
Global reduction	Sustained traffic across a cluster cut	delivered bisection bandwidth; slowest-path completion time	non-blocking capacity, rail alignment, topology-aware collectives
Pipeline point-to-point	Critical-path activation handoffs	startup latency plus size/bandwidth; hop count	short paths, direct placement, transport offload
Expert AllToAll	Many simultaneous worker pairs	path utilization; egress queue depth; tail latency	adaptive routing, path diversity, balanced expert placement
Checkpoint and multi-tenant traffic	Bursts competing with training flows	per-class bandwidth; pause duration; interference	traffic classes (virtual lanes, QoS), admission control

Mapping workload semantics to physical constraints shows that the workload sets the requirements at every level, while each level rests on the one below it. Before higher-level topologies or routing protocols can organize traffic, the physical layer must reliably deliver bits between endpoints without excessive signal attenuation or prohibitive power consumption. Level 1 begins where all fabric architecture originates: at the physical wire.

3.3 Level 1 Wire and Link

At 400 Gb/s and beyond, the physics of signal transmission imposes hard limits on cable length, power consumption, and error rates. Signal attenuation in copper and chromatic dispersion in fiber dictate how far an accelerator can sit from its switches before communication fails. Every higher-level protocol and topology is ultimately constrained by what the physical wire can carry; these limits shape the geometry of the entire cluster.

3.3.1 Signal integrity and PAM4

Achieving 400 Gb/s and higher line rates requires balancing signaling speed against channel loss and detection errors. A port reaches such a rate by summing several lanes, the parallel serial channels that each carry one signal, so a 400 Gb/s NDR port carries four lanes at 100 Gb/s each. The **symbol rate** of a lane counts transmitted symbols per second, each occupying one time slot whether or not the voltage changes. Raising the symbol rate raises the frequencies the channel must carry, and copper loses more signal at higher frequencies. Carrying more bits in each symbol, by using more than two signal levels, raises the bit rate without that penalty, and PAM4 is the scheme high-speed links use to do it.

Definition 3.2: PAM4 signaling

PAM4 signaling is a modulation scheme that represents two bits per symbol with four signal levels, doubling the raw bit rate relative to binary signaling at the same symbol rate.

1. **Significance**: PAM4 increases the raw BW available at a given symbol rate. The reduced separation between levels makes error-control design part of that gain. High-speed Ethernet uses forward error correction (FEC), including Reed-Solomon RS(544,514), with processing latency that depends on the physical-layer (PHY) design and implementation (Anslow 2016; Yin et al. 2022). FEC contributes to the fixed path budget modeled by the α term in section 3.4.4; it is one component alongside propagation, switching, and endpoint processing.
2. **Distinction**: At equal symbol rate and voltage swing, NRZ (Non-Return-to-Zero) binary signaling carries one bit per symbol with greater separation between levels. PAM4 carries two bits per symbol with smaller level spacing; the link must compensate through its signal-quality and error-correction design.
3. **Common pitfall**: Doubling bits per symbol does not guarantee twice the usable payload throughput or a fixed latency penalty. Coding overhead, channel quality, and the chosen PHY implementation determine delivered performance.

The price is noise margin. With four levels packed into the same voltage swing, the gap between adjacent levels in PAM4 shrinks to one third of the NRZ gap (figure 3.4), making the link highly susceptible to noise. Consequently, high-speed links operate near the physical limits of reliable detection and require **forward error correction (FEC)** at the physical layer, commonly using Reed-Solomon RS(544,514)-family codes in high-speed Ethernet designs (Anslow 2016; Ethernet Alliance 2025).

FEC processing is an implementation-dependent latency term; IEEE 802.3df operator material treats roughly 100 ns of host-side FEC with the RS(544,514) code (called KP4 in the standard) and hundreds of nanoseconds of end-to-end module latency as meaningful at AI/HPC fabric timescales (Yin et al. 2022). Because each switch on a path decodes and re-encodes the signal, the FEC cost accumulates with the number of switches a packet traverses.

A path's length is its **hop count**, counted as switch traversals or as inter-switch links; where the two counts differ, each example states which one it uses. A two-tier fabric attaches servers to leaf switches and joins the leaves through spine switches (section 3.5.2), so a packet between leaves traverses three switches. At an assumed 100 to 200 ns per switch, that leaf-spine-leaf route incurs a scenario budget of 300 ns–600 ns of one-way FEC latency, or 600 ns–1200 ns for a round trip.

Against the 350 GB gradient of the 175B job, a few hundred nanoseconds are negligible, because bulk transfers are governed by bandwidth. The cost lands instead on small, latency-bound messages, such as the control handoffs that coordinate pipeline stages, and on collectives whose many sequential steps each pay it. This physical-layer latency sets a floor on the fabric's fixed startup latency, formalized in section 3.4.4 as the α term, that no software optimization can remove.

Doubling line rate without doubling the symbol rate therefore trades signal-to-noise ratio (SNR) for bandwidth, and forward error correction is what lets the link operate that close to the channel limit. Because that tighter margin leaves less room for channel attenuation, it also shortens the distance a passive copper cable can span before the link must switch to optics.

3.3.2 Reach and medium for copper and optics

Past one to three meters, passive copper attenuates high-frequency PAM4 signals beyond what even FEC can recover, forcing any link that exits a rack onto active optical cables or pluggable optics at a steep jump in cost and power. The choice of physical medium therefore determines how far a link can reach before latency, power, and cabling budgets reshape the cluster geometry, and the recurring cable and optics options differ along exactly those attributes (table 3.3).

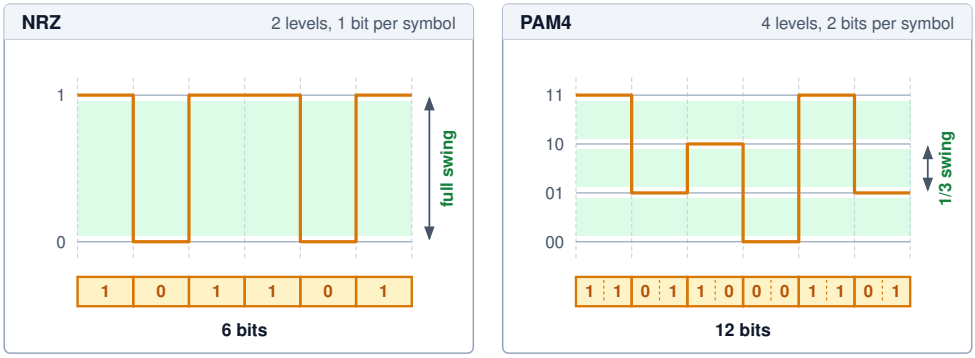

Figure 3.4: From NRZ to PAM4 Signaling: Both panels send six symbols in the same time over the same total voltage swing. NRZ (left) uses two levels and carries one bit per symbol, 6 bits in all; PAM4 (right) uses four levels and carries two bits per symbol, 12 bits in all, doubling the data rate without raising the symbol rate. The cost is noise margin: adjacent PAM4 levels sit one third as far apart (arrows), so high-speed links need forward error correction (FEC), which adds to the fabric's fixed α latency.

Table 3.3: Cable and Optics Placement Trade-Offs: Common interconnect media occupy different reach, latency, and cost points, so cluster designers use copper inside racks and optics as distance begins to dominate placement.

Medium	Reach	Cost profile	Latency impact	Typical placement
DAC (direct attach copper)	1 to 3 meters, shrinking as lane rate rises	Lowest cost (~$50) and no optical module power	Lowest media latency; no optical conversion, though high-speed lanes still pay FEC	Within a rack
AOC (active optical cable)	3 to 30 meters	Higher cost (~$500) from permanently attached transceivers (the modules that convert between electrical and optical signals)	Electrical/optical conversion adds latency on top of the forward error correction (FEC)	Rack-to-rack or row-scale links
Pluggable optics	More than 30 meters	Highest module and power cost, but replaceable optics and fiber cords	Optical conversion adds latency on top of FEC; distance enables topology reach	Network core and pod-scale links

The high-speed copper and optical paths considered here both include FEC, so choosing optics does not by itself remove error-correction overhead. The ranges in table 3.3 describe representative placements rather than fixed boundaries between media. Optical interconnects can extend the layout beyond a passive copper channel's budget, while pluggable optics also permit separate replacement of modules and fiber. Reach, delivered latency, power, and serviceability jointly determine the choice.

Systems Perspective 3.2: The cost of distance

In an ML fleet, distance is money. Assuming two spine-layer optical links per accelerator, a cluster of 10,000 GPUs requires ~20,000 optical links at the spine layer alone. At $500 per complete link, that represents $10 million in cabling. For power, count one selected spine-side pluggable module per link at 20 W per module, giving 400 kW for those endpoints alone. Each optical link also has a module at its opposite end; this subtotal excludes those modules, switch electronics, and cooling, so it is not the full fabric power budget. The cost dictates cluster geometry: architects pack accelerators as densely as possible (70–100 kW per rack) to maximize cheap copper and minimize expensive optics.

3.3.3 SerDes, link budget, and power

Driving thousands of high-speed links across a cluster burns tens of kilowatts just to serialize bits onto the wire, a noticeable share of the accelerators' own thermal budget. Every port depends on a **serializer/deserializer (SerDes)** circuit that converts parallel data from the switch application-specific integrated circuit (ASIC) into a serial stream, and its power cost accumulates across the provisioned ports. For the 1,024-GPU reference cluster, assume a full-BW_{bisect} fabric with roughly 3,000 high-speed links across its host and switch tiers. Count one selected switch-side port per link in this power subtotal. Each counted 400 Gb/s port is assigned 25 W, more than the 20 W pluggable-module figure in the cost-of-distance estimate because it also counts the switch-side SerDes that drives the module. The model totals 75 kW, equivalent to 10.5 percent of the accelerators' combined thermal design power. This comparison isolates the assumed power of the selected switch-side ports, excluding the opposite endpoints and the rest of the switch chassis; it does not establish total fabric power or the fraction of measured cluster energy consumed by networking.

The **link budget** allocates the signal loss a path can tolerate while meeting its error target. Losses are counted in decibels (dB), a logarithmic power ratio in which every 10 dB is a further tenfold attenuation (3 dB is roughly a halving), so the losses of cables and connectors in series add directly. In this simplified accounting, a lane has a 30 dB allowance. A 3-meter DAC cable contributing 18 dB of loss and connectors contributing 2 dB leave 10 dB of margin. Higher signaling frequencies increase copper loss, so preserving reach can require more signal conditioning, a different cable, or optics. Aggregate port speed alone does not determine cable reach; lane rate, channel construction, and active components matter. The design must meet both the signal-quality target and the power budget before the transport layer can rely on the link.

3.4 Level 2 Transport and the Performance Model

Closing the physical link budget delivers raw bits between adjacent ports, but a single AllReduce across 1,024 GPUs still stalls if host software sits in the path of every transfer or if a dropped packet triggers a multi-millisecond recovery pause. Turning raw wire rate into predictable step time requires a transport layer that offloads memory-to-memory delivery and bounds loss recovery (a contract met either through a physically lossless fabric or through loss-tolerant endpoint recovery), together with a performance model that separates fixed per-message startup latency from sustained link bandwidth.

3.4.1 RDMA and GPUDirect

A conventional kernel TCP path charges the host for every byte it moves. Interrupt processing, protocol work, and payload copies consume CPU cycles and memory bandwidth that the training process also needs for data loading and preprocessing, and the charge grows with aggregate I/O as many streams or NICs drive a host at 400 Gb/s link rates. Modern TCP stacks reduce the charge with segmentation offload, receive-side scaling, zero-copy interfaces, and careful non-uniform memory access (NUMA) placement, so a single stream need not exhaust a server's memory bandwidth. They do not remove it, and the contention becomes harder to provision predictably as nodes add NICs.

Definition 3.3: Remote direct memory access (RDMA)

Remote direct memory access (RDMA) is a networking technology that offloads transport processing to network interface hardware, allowing access to registered remote memory without per-transfer kernel processing of the payload.

1. **Significance**: RDMA can reduce end-to-end message latency from tens of microseconds for a kernel TCP path to low-single-digit microseconds on a tuned fabric, shrinking the fixed startup cost that every message adds to the Fleet Law's communication term $T_{\text{comm}}(N)$.
2. **Distinction**: Unlike a conventional kernel TCP path, RDMA offloads transport processing to NIC hardware and permits user-space access to registered memory. The resulting CPU-core and memory-bandwidth savings depend on the host, NIC, I/O path, message size, and software configuration rather than on link rate alone.
3. **Common pitfall**: RDMA is an access and offload model, not one loss-recovery algorithm. Reliable InfiniBand and RoCE transports detect loss and retransmit (InfiniBand Trade Association 2000). Some RNIC generations use the Go-Back-N-style recovery noted with the loss inversion, which makes a drop expensive (section 3.4.3 quantifies the cost), while other designs support stronger selective recovery or tolerate lossy Ethernet without the link-layer pause mechanism (priority flow control) that many RoCE deployments rely on (Mittal et al. 2018). The requirement is reliable completion with a recovery tail that the synchronized workload can tolerate, not a physically lossless Ethernet fabric in every implementation.

RDMA bypasses the kernel data path. Before any transfer, software registers a memory buffer, which pins its pages in physical memory and gives the NIC the addresses and access rights it needs. The sending NIC can then name a registered buffer on the remote machine, and the remote NIC places the data there directly, with no remote CPU copying the payload. Offloading transport logic to the NIC in this way allows applications to read and write registered remote memory without per-message kernel intervention. Each message skips the interrupts, software protocol work, and payload copies of the kernel path, which is how RDMA shrinks the fixed startup cost.

For ML, GPUDirect RDMA,[3] the path that dense nodes pair with one NIC per GPU (section 2.5.2), extends this principle to accelerator memory (NVIDIA 2026b). Within a sending node (figure 3.5), GPUDirect RDMA routes an n-byte payload (what one participant sends, not the model-wide gradient state) along the solid green peer-to-peer path through the PCIe switch straight to the RDMA NIC, bypassing host memory entirely. Without this direct memory access (DMA) bypass, the dashed gray staging path must first copy n bytes from GPU memory into a registered host-memory buffer and subsequently read n bytes out to the NIC, generating roughly $2n$ bytes of redundant host-memory bus traffic. The network still transfers the payload, and the PCIe path may comprise multiple DMA transactions; the optimization removes staging copies rather than collapsing the transfer into one transaction.

[3] **GPUDirect RDMA**: Introduced by NVIDIA for Kepler-class GPUs and CUDA 5.0, GPUDirect RDMA enables a direct PCIe path between GPU memory and third-party peer devices such as network interfaces (NVIDIA 2026b). Before GPUDirect, every gradient transfer had to stage through host memory, consuming CPU memory bandwidth that competes with data loading. Eliminating this bounce path is one reason overlapping communication with backward-pass computation is feasible at scale.

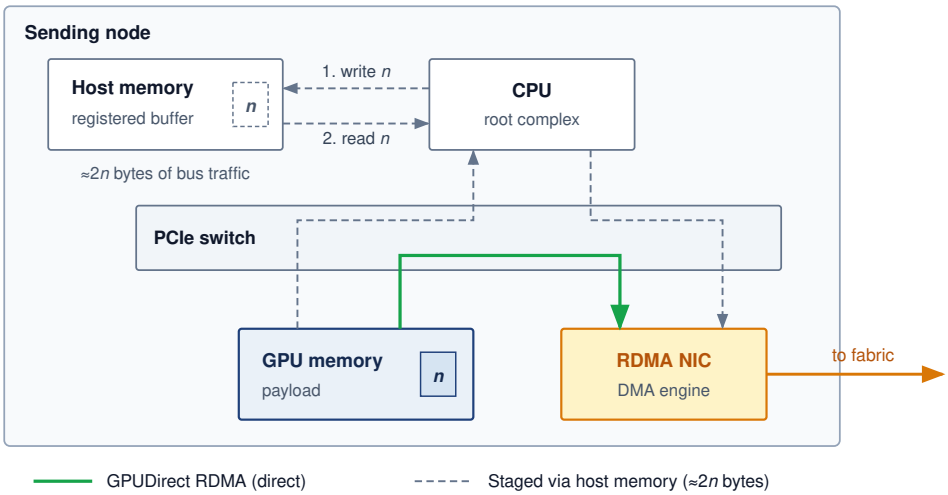

Figure 3.5: GPUDirect RDMA Data Path: Without GPUDirect (dashed), a sending node first writes the n-byte payload from GPU memory into a registered host-memory buffer and the NIC then reads it back out, about $2n$ bytes of host-memory traffic per payload. With GPUDirect RDMA (solid), the NIC reads GPU memory directly over the PCIe peer-to-peer path, so the host-memory bounce disappears.

3.4.2 InfiniBand and RoCE

GPU clusters commonly choose between two RDMA transport stacks, and the decision is a reliability-versus-operations trade-off. In their usual deployments, both stacks keep switch buffers from overflowing under synchronized bursts, and they do so in different places. InfiniBand uses **credit-based flow control**. Each receiver advertises its free buffer space to the sender at the other end of the link in units called credits, and the sender transmits only into space that has been advertised. A buffer therefore cannot overflow; congestion appears as senders waiting for credits rather than as dropped packets. Ethernet was designed to drop packets when buffers fill, so RoCE deployments that want the same behavior add two mechanisms. Priority flow control (PFC) lets a filling switch port send a PAUSE frame that stops its upstream neighbor for one traffic class, and explicit congestion notification (ECN) marks packets that pass through a filling queue so that their senders slow down before a pause is needed. InfiniBand thus builds losslessness into the link layer, while RoCE assembles it from switch configuration and end-to-end congestion control (table 3.4); section 3.6.1 develops PFC's failure modes.

Table 3.4: RDMA Transport Stack Trade-Offs: InfiniBand and RoCE expose RDMA semantics through different operational contracts, one built around native losslessness and the other around Ethernet compatibility.

Stack	Fabric behavior	Operational burden	Typical fit
InfiniBand (IB)	Purpose-built HPC switched fabric (InfiniBand Trade Association 2000) with credit-based flow control, so losslessness is native at the link layer	a central subnet manager computes routes, and virtual lanes (VLs), separate buffer and credit pools on one link, isolate traffic	Dedicated training clusters where predictable tail latency outweighs Ethernet ecosystem flexibility
RoCE (RDMA over Converged Ethernet)	RoCEv2 carries RDMA semantics by encapsulating InfiniBand transport headers in UDP/IP packets	PFC-based deployments suppress loss with PFC, ECN, and routing policy (Guo et al. 2016; Gangidi et al. 2024); PFC-free designs instead place more burden on end-to-end congestion control and loss recovery (Mittal et al. 2018)	Multi-vendor Ethernet fleets that can absorb more congestion-control tuning

[4] **InfiniBand**: Formed in 1999 from the merger of two competing server I/O standards (Future I/O and NGIO), InfiniBand was originally designed to replace PCI as a general-purpose system interconnect. Its pivot to HPC networking preserved the credit-based, hardware-managed flow control that server I/O demanded, and that heritage is why InfiniBand provides native losslessness without the PFC fragility of Ethernet-based RDMA fabrics.

The InfiniBand[4] row reflects a protocol heritage that favors hardware-managed losslessness over Ethernet compatibility. Both RDMA stacks nonetheless share the same user-space bypass across the user-space, kernel, and hardware boundary (figure 3.6). The conventional copied TCP/IP path on the left uses two software copies across kernel buffer boundaries and runs the protocol in software. The RDMA stack on the right traces a single amber bypass arrow from the collective communication library (such as the NVIDIA Collective Communications Library, or NCCL) and the user-space **verbs API** straight to the RDMA NIC, which keeps the operating system off the active data path. The verbs API is the programming interface through which software posts RDMA work requests to the NIC.

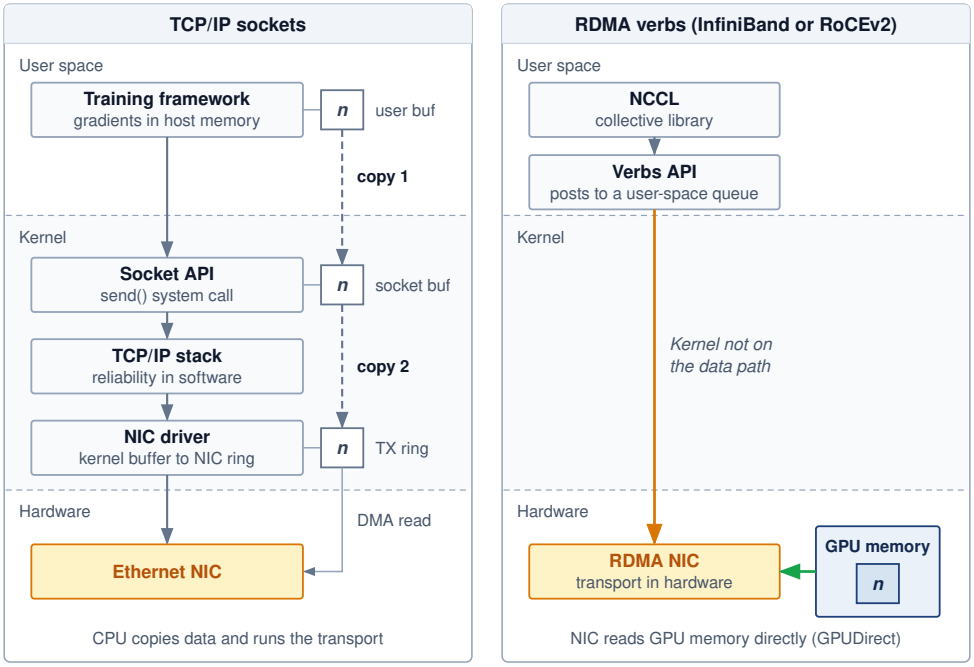

Figure 3.6: Traditional TCP/IP vs. RDMA Verbs Stacks: Read each stack top to bottom through the user-space, kernel, and hardware bands. In the illustrated copied TCP/IP socket path (left), a send enters the kernel through a system call, the payload is copied into kernel socket buffers and again toward the NIC's transmit ring, and the CPU runs the transport protocol in software. With RDMA verbs over InfiniBand or RoCEv2 (right), the collective library posts work requests from user space, the kernel stays off the data path, and the RDMA NIC runs the reliable transport in hardware, reading GPU memory directly through GPUDirect instead of staging a host copy.

3.4.3 Reliable delivery and loss recovery

By routing the user-space verbs API directly to hardware (figure 3.6), the architecture decouples collective library calls from the operating system, leaving the transport hardware and fabric to arbitrate reliable delivery and packet loss. Reliable connected transports in both InfiniBand and RoCE detect loss and retransmit (InfiniBand Trade Association 2000); credits and PFC exist to make that retransmission rare (Guo et al. 2016; Gangidi et al. 2024). Losslessness is therefore a deployment contract rather than a property of RDMA as a family. Designs that operate over lossy Ethernet without PFC shift more responsibility to endpoint congestion control, packet reordering, and selective recovery (Mittal et al. 2018).

Recovery granularity therefore matters. On an RNIC with **Go-Back-N**-style recovery and no out-of-order retention, a lost packet invalidates every packet that arrives after it, and the sender must retransmit from the missing sequence point. The penalty is everything the sender transmitted between the lost packet and the moment it learned of the loss. That span is the resent window,

and it is at most link rate times detection delay. If the receiver's negative acknowledgment (its message reporting the gap) reaches the sender after one 5 µs round trip, a sender at 50 GB/s has meanwhile sent 250 KB, all of which it must send again; a drop 900 MB into a 1 GB transfer wastes that window, not the 900 MB already delivered. The transfer stalls until the resent window arrives. When detection instead waits for the sender's own retransmission timeout, the stall grows with that timeout, although the resent window stops growing once the sender reaches its cap on data sent but not yet acknowledged. Resending at line rate takes at most one more detection delay, because the window is at most link rate times that delay. Below that cap, a faster link therefore resends more bytes per drop without shortening the stall.

An RNIC with selective recovery retransmits only the missing packets, so the same physical loss produces a shorter tail (Mittal et al. 2018). Either way the stall spans many packet wire times, and a synchronous collective exposes the slowest recovery to every waiting worker.

Fabric designers can reduce that tail by suppressing drops with PFC, by using a loss-tolerant transport with bounded endpoint recovery, or by combining congestion control and admission policy so queues rarely overflow. Each choice moves state and failure modes between switches and endpoints.

3.4.4 The α-β performance model

Even on a lossless RDMA path with zero packet drops, doubling the link rate barely speeds up a few-kilobyte pipeline control handoff while cutting a multi-hundred-megabyte gradient transfer nearly in half. Separating those two behaviors requires a cost model that distinguishes fixed per-message overhead from per-byte transmission rate. The α-β model (principle 5) decomposes message transfer time as $T(n) = \alpha + n/\beta$, where n is the message size in bytes, α is the fixed startup latency, and β is the sustained bandwidth in bytes per unit time (Hockney 1994). The physical topology influences both parameters through path length and available capacity. A collective built from many sequential rounds of messages pays α once per round, so its startup cost can grow with the number of participants even when the underlying fabric provides short physical paths.

The model reveals two regimes that call for different engineering responses. When $n < \alpha\beta$, the startup cost α exceeds the transmission cost, placing small control messages deep in the latency-bound region, whereas when $n > \alpha\beta$, the n/β term dominates and large gradient transfers sit deep in the bandwidth-bound region. Near the crossover both terms matter, and where a collective lands also depends on how its algorithm divides the payload into messages and schedules them across rounds.

↪ The same decomposition, summed over each schedule's rounds, is developed for collective algorithms in section 5.5.

For the modeled NDR InfiniBand path with $\alpha = 1.5$ µs and $\beta = 50$ GB/s, the crossover point $n^* = \alpha \cdot \beta \approx 75$ KB. Measured startup latency and delivered bandwidth vary with the NIC, switch path, software stack, and contention, so these parameters describe a scenario rather than a protocol constant. Far below the crossover, more bandwidth removes only a small part of transfer time; far above it, lower startup latency removes only a small part. Shorter paths can lower α, while additional delivered path bandwidth raises β; their value depends on the message sizes the workload actually exchanges.

↪ The crossover message size that separates the latency-dominated and bandwidth-dominated regimes, with a worked InfiniBand NDR example, is developed in section D.1.1.

Napkin Math 3.1: The alpha-beta crossover

Problem: A fabric designer is comparing a modeled InfiniBand NDR path (1.5 µs, 50 GB/s) against a modeled RoCE path over slower Ethernet (5 µs, 12.5 GB/s). The comparison deliberately bundles a higher-latency stack with a slower link, so a third path, RoCE on a 400 Gb/s link, separates the two effects. For a 4 KB control message and a 350 MB gradient shard, which part of $T(n) = \alpha + n/\beta$ dominates, and why does the first path help for different reasons in each regime?

Math: Apply $T(n) = \alpha + n/\beta$ for a 4 KB control message and a 350 MB gradient shard.

1. **Small message** (4 KB):
   - InfiniBand: $1.5\,\mu s + 4\,KB/50\,GB/s = 1.5 + 0.08 = \mathbf{1.58}\,\mu s$
   - RoCE (100G): $5.0\,\mu s + 4\,KB/12.5\,GB/s = 5.0 + 0.32 = \mathbf{5.32}\,\mu s$

- **Result**: Under these assumed path parameters, the NDR path is 3.4× faster because its modeled α is lower.

2. **Large message** (350 MB):
   - InfiniBand: $T = \alpha + n/\beta = 7.0$ ms
   - RoCE (100G): $T = \alpha + n/\beta = 28.01$ ms
   - **Result**: Under these assumed path parameters, the NDR path is 4× faster because its modeled β is higher. That factor is the ratio of the two line rates, so it measures the faster link rather than the transport.
3. **Same line rate** (RoCE at 400 Gb/s, RoCE α):
   - Small message: 5.08 µs, still 3.2× the NDR time, because the assumed α is all that separates these two paths.
   - Large message: 7.0 ms, essentially the NDR time, because at equal β nothing but a few microseconds of α separates them.

Systems insight: This scenario places the crossover ($n^* = \alpha\beta$) near 75 KB. Because gradients are megabytes to gigabytes, their transfers are usually bandwidth dominated, whereas pipeline coordination can remain latency dominated. A higher line rate produces little benefit for a message whose startup term dominates; transport processing, topology, and hop count then govern the improvement. At equal line rate the two stacks in this scenario differ on the control message and barely at all on the gradient shard, so a recommendation between them weighs how much step time small messages consume against the operational burden of each stack (table 3.4).

The two message sizes in notebook 3.1 map directly onto the 175B training job. For the 4 KB control message that notifies a pipeline stage its next microbatch is ready ($T(n) \approx 1.58$ µs on NDR), the transmission term contributes only 5.1 percent of the total, so faster links offer little benefit. For the rounded 350 MB shard that stands for one share of the 350 GB gradient across the 1,024 GPUs ($T(n) \approx 7.0$ ms for that one message), transmission dominates, and doubling bandwidth to 100 GB/s approximately halves the transfer time.

Across the full range of message sizes, each path traces a flat latency-dominated plateau that bends into a linear bandwidth-dominated regime (figure 3.7). On the modeled NDR path, the knee at 75 KB marks where transmission time overtakes the 1.5 µs startup latency; the 100 Gb/s RoCE path sits on a higher floor of 5 µs and bends near 62.5 KB. The exact positions reflect the assumed α and β scenario parameters rather than invariant protocol constants. On both paths, startup latency governs sub-kilobyte control handoffs, whereas multi-megabyte gradient shards sit well inside the bandwidth-bound regime.

Beyond classifying individual transfers, the same model identifies when a collective overtakes computation. A ring AllReduce, one common schedule, passes chunks around a logical cycle of participants, in a ReduceScatter phase that leaves each participant one chunk of the combined result and then an AllGather phase, with no requirement that the physical switches form a ring. The scaling laws that push training toward more parameters (section 1.3) enlarge, in direct proportion, the gradient that every data-parallel step must reduce. Summed over the ring's rounds, the α-β cost of each transfer then decides when that collective outlasts the step's computation. At that point the network becomes the pacing constraint for the entire job.

↪ The ring AllReduce derivation, and the schedules that compete with it, are developed in section 5.6.3.

A measured collective shows both regimes and where the two-parameter model needs care. NVIDIA publishes the full output of the standard NCCL AllReduce benchmark on one node of eight H100 GPUs connected by NVLink, sweeping message sizes from 8 bytes to 16 GB (NVIDIA 2026a). The run crosses no external network, so it tests the two-parameter form itself, free of external switch paths and fabric congestion. The two-parameter model can be built from two summary values of that sweep, the median time of the small messages as α and the delivered rate of the largest message as β, and set beside the measured time (figure 3.8). The small messages sit on a floor of about 17 µs

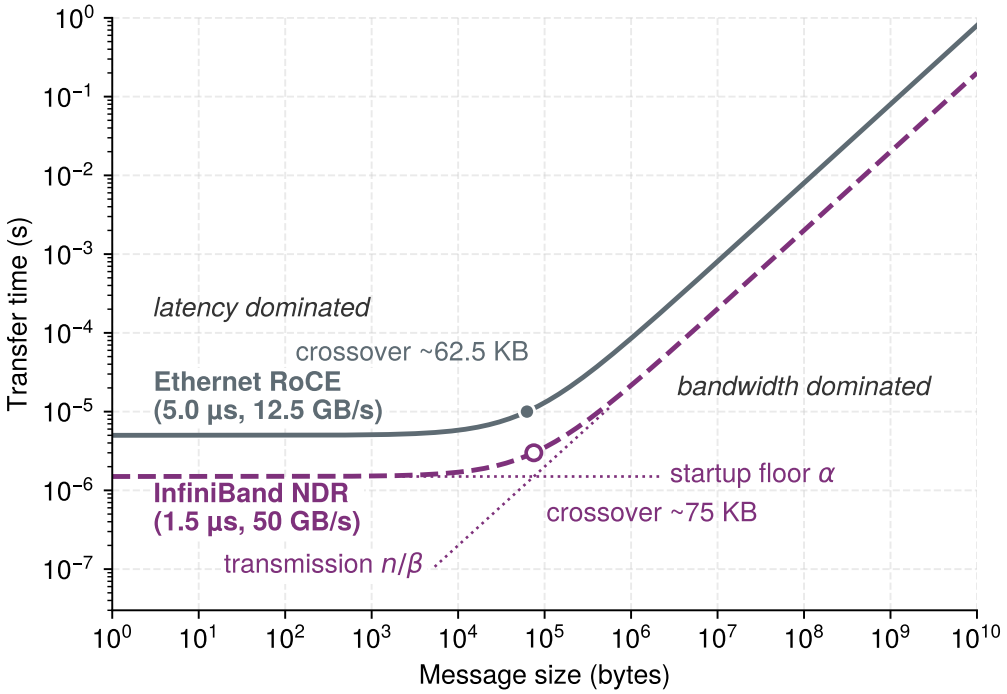

Figure 3.7: The α-β Crossover: Modeled transfer time as a function of message size for an NDR InfiniBand path and a 100 Gb/s RoCE baseline using the scenario parameters in the text. Small messages are latency dominated (flat region), while large messages are bandwidth dominated (linear region); on the InfiniBand path the dotted asymptotes show the knee sitting where the two terms are equal. Measured crossover points vary with the end-to-end path.

regardless of size, and the largest messages approach 274 GB/s of algorithm bandwidth, defined as the tensor size divided by completion time.

The benchmark multiplies that algorithm bandwidth by $2(8-1)/8 = 1.75$, the bytes a ring AllReduce sends across each GPU's links per tensor byte reduced, to report 480 GB/s of bus bandwidth. The benchmark applies this factor whichever schedule ran, so the figure is a normalization rather than a link measurement, and it can exceed NVLink's 450 GB/s per direction.

Between the two extremes, the model places this fitted collective's knee near 5 MB, far above the 75 KB crossover of a single NDR path, but the measured time runs up to 1.7× the model's prediction for messages between roughly 1 MB and 1 GB. Across that middle range, chunks are not yet large enough to fill every parallel NVLink channel and amortize per-round protocol overhead, so delivered bandwidth is still ramping toward its peak. The model is therefore close at the two ends and optimistic in the middle, and that middle range is where the gradient buckets (the groups of gradient tensors reduced in one call) of data-parallel training often fall.

↪ In-network reduction, in which switches rather than GPUs combine contributions and which NVLink SHARP brings to the node's NVLink switches, is developed in section 5.7.3.

Checkpoint 3.1: Protocol selection

Consider the 175B job on the 1,024-GPU reference cluster. Its gradient exchanges move messages of hundreds of megabytes, while its pipeline stages coordinate through frequent control messages of a few kilobytes.

- ☐ **Protocol fit**: use the α-β crossover to recommend a transport stack, and justify the choice for both the large gradient messages and the small control messages.
- ☐ **RoCE requirements**: compare the infrastructure required by PFC-based and PFC-free RoCE designs with InfiniBand's credit-based fabric.

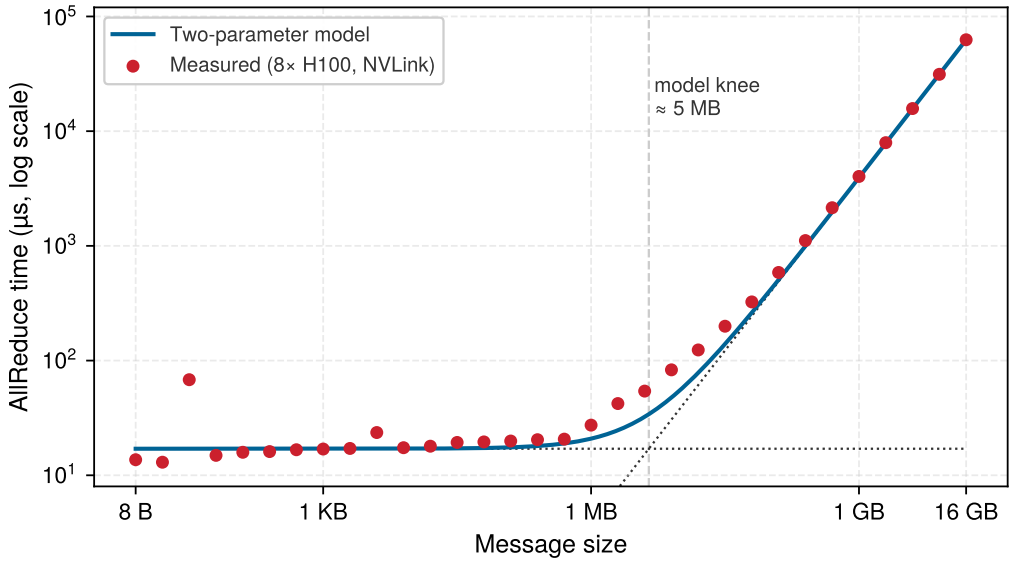

Figure 3.8: Measured AllReduce Time against the Two-Parameter Model: NCCL AllReduce time on one node of eight H100 GPUs over NVLink for message sizes from 8 bytes to 16 GB, on logarithmic axes (points), with the two-parameter model built from the measured small-message floor and the largest message's delivered bandwidth (solid line) and its two asymptotes (dotted). The model is close at both ends; between roughly 1 MB and 1 GB the measurement runs above it. One early small-message point is a timing outlier. Data from NVIDIA (2026a); whether NVLink SHARP was enabled for this run is not stated.

☐ **Line rate vs. startup cost**: explain why moving the control messages from a 100 Gb/s path to a 400 Gb/s path with the same α would barely change their transfer time.

The α-β model quantifies the speed of a single link, but the 175B-parameter model requires 1,024 GPUs to work in concert. Scaling these transport primitives from a pair of nodes to a warehouse-scale supercomputer demands a **network topology**: a specific pattern of connections that maximizes $\text{BW}_{\text{bisect}}$ while minimizing the hop count and cabling cost for global collective operations.

3.5 Level 3 Switch and Topology

A cluster can have thousands of fast links at the edge and still starve its AllReduce operations if the cross-sectional capacity at the narrowest point in the switching hierarchy is insufficient. The physical arrangement of switches determines whether that cross-section throttles collective communication or sustains full line-rate exchange across all participating workers. Two properties govern this capability: bisection bandwidth $\text{BW}_{\text{bisect}}$ and whether the fabric is non-blocking. The bisection bandwidth theorem (principle 6) captures the constraint, stating that traffic crossing a cut is limited by the aggregate bandwidth of that cut. Bisection bandwidth quantifies the worst-case throughput ceiling that the topology imposes on global collectives.

Definition 3.4: Bisection bandwidth

Bisection bandwidth is a network topology metric defined as the minimum aggregate link capacity crossing any partition that divides the cluster into two equal halves, representing the worst-case throughput ceiling for traffic that must cross that partition, including global collectives.

1. **Significance**: Bisection bandwidth $\text{BW}_{\text{bisect}}$ directly sets the cluster-scale bandwidth ceiling for global synchronization. A 1,024-GPU fat-tree with 400 Gb/s links at 1:1

subscription provides $BW_{bisect} = 512 \times 50\,GB/s = 25.6\,TB/s$ per direction; a 4:1 oversubscribed spine reduces this to 6.4 TB/s, increasing the transfer-time lower bound for a fixed cut-crossing payload by 4×. Actual AllReduce slowdown depends on rank placement, the collective schedule, and which links are saturated.

2. **Distinction**: Unlike aggregate bandwidth (the sum of all edge link speeds, which can be high even in a poorly connected topology), BW_{bisect} measures global connectivity. A star topology with 1,000 edge links all meeting at one central switch has high aggregate bandwidth but BW_{bisect} limited by that switch's backplane capacity.
3. **Common pitfall**: A frequent misconception is that adding more leaf switches always increases BW_{bisect}. In a three-tier fat-tree, oversubscribing the spine layer (giving each spine switch fewer uplinks to the core than downlinks to its leaves) reduces BW_{bisect} below the edge-link total regardless of how many leaf switches are present.

ML training includes global collectives and AllToAll phases that can drive traffic across topology partitions, making full BW_{bisect} a useful capacity target. A smaller cut becomes a bottleneck when the workload's crossing traffic exceeds what it can deliver within the communication budget, and rank placement together with the collective schedule determines how much traffic crosses that cut.

Bisection bandwidth measures cut capacity; a non-blocking fabric provides paths that can sustain simultaneous input-output pairings at line rate with suitable routing. Full BW_{bisect} and balanced uplink capacity support that goal, but routing must also spread traffic across the available paths. In the oversubscribed designs considered here, BW_{bisect} shrinks by a factor equal to the oversubscription ratio. The transfer-time bound for a fixed payload that saturates the cut rises by that same factor, whereas a collective that keeps most traffic within each half need not experience the same slowdown.

↪ A slower cut lengthens the training step only when the extra transfer time cannot hide behind computation. The communication-to-compute diagnostic that tests this is developed in section B.4.1.

Definition 3.5: Non-blocking fabric

Non-blocking fabric is a network with sufficient internal path capacity to support any permutation of input-output port pairs simultaneously at full line rate under a suitable routing assignment.

1. **Significance**: In ML fleets, a non-blocking fabric can support full-rate permutation traffic with suitable routing, preserving the available BW_{bisect} term for global collectives. A 2:1 oversubscribed spine, a common data-center compromise, halves the effective BW_{bisect}, doubling the communication time when that cut is the saturated bottleneck and dropping scaling efficiency $\eta_{scaling}$ accordingly. In a scenario workload assumed to spend 30 percent of each step communicating, doubling that share while the computation share stays fixed lengthens each step, and the longer step costs roughly 23.1 percent of total cluster throughput. The bisection definition above and notebook 3.2 use 4:1 instead, a cost-optimized spine; the notebook writes out the throughput formula behind both results.
2. **Distinction**: The guarantee concerns a permutation, in which each input and output belongs to one pairing. It does not promise full line rate from every sender to every receiver at once, nor does it make an arbitrary routing assignment free of contention.
3. **Common pitfall**: A frequent misconception is that non-blocking means zero congestion. Endpoint congestion (incast) can still occur if multiple senders simultaneously target the same receiver port, regardless of how much internal fabric capacity is available.

Fat-trees build on Clos-style non-blocking network principles to supply full BW_{bisect} (Clos 1953), rail-optimized networks trade some global routing flexibility for lower cabling cost, and dragonfly topologies reduce optical link counts while introducing workload-placement constraints (Kim et al. 2008). This trade-off between guaranteed bisection bandwidth and cabling cost drives every topology decision in large-scale ML clusters, as the comparisons in section 3.5.2, section 3.5.3, and section 3.5.4 show, and every switched design begins at the first-hop switch that aggregates a rack or rail of servers.

3.5.1 Top-of-rack (ToR) and the failure domain

Consider a rack containing four nodes with eight GPUs each: if all 32 accelerators attach to a single top-of-rack (ToR) switch with no alternate path, losing that switch disconnects the entire rack from the job at once. Because the ToR aggregates every server link in the rack, it becomes a shared failure domain, one level of the physical hierarchy defined in section 2.5.5, and recovery then depends on available paths, surviving state, and the job's recovery policy. This one-switch-per-rack wiring is the simpler baseline used here to teach radix and failure domains. The reference cluster itself wires each GPU to a rail switch, a layout developed in section 3.5.3. That layout reshapes the failure domain, because losing one rail switch cuts one GPU in each of many nodes rather than every GPU in one rack.

The same boundary shapes state placement. To survive a rack-wide outage, a job needs recoverable state outside that rack and a way to resume on available capacity. Topology-aware scheduling must therefore consider both fast communication paths and the failures that several nearby resources share.

↪ Nested failure domains, and why redundant state must span them to survive a correlated rack or ToR outage, are developed in section 7.3.4.

The switch **radix**,[5] its number of ports, determines how many endpoints and uplinks it can connect. In the reference allocation, a 64-port switch assigns 32 ports to servers, one for each of 32 GPUs, and 32 ports to the spine. With equal port rates, the 1:1 ratio avoids oversubscription at that boundary; delivered throughput still depends on forwarding and the paths beyond it. The 1,024-GPU cluster therefore requires approximately 32 such switches, although the physical rack count also depends on node density and the cabling layout.

[5] **Switch radix**: High-end switches such as NVIDIA Quantum-2 offer a radix of 64 ports, each at 400 Gb/s. As model scale pushes toward 100,000 GPUs, higher radix (128 or 256 ports), larger local accelerator domains, or accepted topology constraints are the main ways to avoid adding switch tiers, each of which adds switch latency and cost.

3.5.2 Fat-tree (Clos) networks

The ToR provides non-blocking bandwidth within a single rack, but connecting the reference cluster's 32 first-hop switches into a cluster that preserves full $\text{BW}_{\text{bisect}}$ across every possible communication pair requires a topology that scales cross-sectional capacity with cluster size. The fat-tree (also called a Clos network) achieves this by adding parallel switches at each tier above the leaves, so the aggregate uplink capacity always matches the total edge bandwidth feeding into it. In this vocabulary, leaf switches attach servers or racks, spine switches connect leaves within a pod, and core switches connect multiple pods when a third tier is needed; other sources call the leaf and spine tiers edge and aggregation. A pod here is a switching domain, the group of leaves that one set of spines joins, which can be narrower than the multi-rack pod of section 2.7. A multi-rack pod in that sense can contain several fat-tree pods.

Definition 3.6: Fat-tree

Fat-tree is a hierarchical ML cluster network topology in which the number of parallel paths (and therefore aggregate cross-sectional capacity) increases at each switch tier toward the spine, providing full $\text{BW}_{\text{bisect}}$ and multiple equal-cost routes between any two nodes (Al-Fares et al. 2008).

1. **Significance**: A k-ary fat-tree built from radix-k switches supports $k^2/2$ hosts in a standalone two-tier configuration and $k^3/4$ hosts in a three-tier configuration with full $\text{BW}_{\text{bisect}}$. With 64-port switches, a standalone two-tier fabric supports 2,048 GPUs and a three-tier fabric supports 65,536 hosts. This is a capacity bound for suitable routing and admissible traffic; it does not eliminate endpoint contention or guarantee that a particular AllReduce implementation reaches line rate.
2. **Distinction**: Unlike a standard tree (where bandwidth at the root is a single bottleneck shared by all leaves), a fat-tree replaces each root with multiple spine switches whose combined uplink capacity matches the total edge bandwidth, eliminating the bottleneck.
3. **Common pitfall**: A frequent misconception is that fat-trees guarantee zero network cost. They require $\mathcal{O}(N \log N)$ switches and dense cabling, so a non-blocking three-tier fat-tree at the 4,000-GPU scale needs hundreds of switches and more than ten thousand links before a single accelerator is purchased.

A three-tier radix-4 fabric ($k = 4$) shows the non-blocking structure at a scale small enough to draw (figure 3.9), with four core switches interconnecting four dashed pod modules. Leaf and spine switches split their ports between the tiers below and above; core switches use all their ports to connect pods. Each tier boundary therefore carries 16 links. The highlighted route traverses five switches and is one of $(k/2)^2 = 4$ equal-cost routes from a GPU in Pod 1 to Pod 4. Parallel paths supply the capacity, while routing must distribute traffic across them to use it.

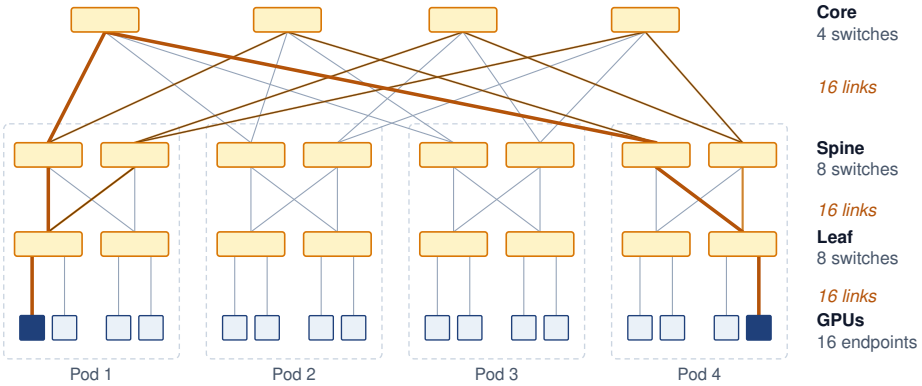

Figure 3.9: Non-blocking Fat-Tree Topology: A three-tier fat-tree built from radix-4 switches ($k = 4$) connects $k^3/4 = 16$ GPUs through four pods. Leaf and spine switches split their ports equally between the tiers below and above; core switches use all their ports to connect pods. Each tier boundary carries 16 links, supporting full $\text{BW}_{\text{bisect}}$ with suitable routing. The highlighted leaf-spine-core-spine-leaf route between pod 1 and pod 4 is one of $(k/2)^2 = 4$ equal-cost paths (the other three are drawn lighter), and exploiting those alternatives without excessive reordering requires compatible switch and endpoint mechanisms.

The fat-tree[6] is a common default for ML clusters because a non-blocking design can provide full $\text{BW}_{\text{bisect}}$ across arbitrary partitions. This capacity benefits global AllReduce even though implementations such as rings and trees do not require every pair of ranks to communicate at once. The number of hosts that the guarantee can cover depends on switch radix and tier count.

A fat-tree built from switches with radix k supports $N_{\text{hosts}} = k^3/4$ hosts at three tiers. A standalone two-tier non-blocking fabric can instead use k leaves, each with $k/2$ host ports and $k/2$ uplinks, connected to $k/2$ spines. Each spine uses all k ports to reach the leaves, providing $k^2/2 = 2{,}048$ host ports for radix-64 switches. In the three-tier design, each spine switch reserves half of its ports for the core tier and uses the other $k/2$ ports to connect to leaves, so one pod has $k/2$ leaves and $(k/2)^2 = 1{,}024$ host ports. Because each radix-k core switch connects one port to each pod, the core tier can join up to k pods for the full $k \times (k/2)^2 = k^3/4$ host capacity. These are different port allocations and switch counts rather than two ways of counting the same fabric, and while the 1,024-GPU reference cluster fits within a standalone two-tier design, the distinction matters as soon as a cluster outgrows that limit.

A 4,096-GPU deployment cannot join two fully populated standalone 2,048-host fabrics at their spines, because every spine port is already occupied by a leaf connection. Instead, it is built as four 1,024-host pods of the three-tier design, joined through a core tier. A concrete bisection estimate (notebook 3.2) shows how quickly oversubscription at that spine or core boundary turns into synchronization delay.

6 **Fat-tree (Clos network)**: The underlying multi-stage switching theory was invented by Clos (1953) at Bell Labs to minimize the number of electromechanical crosspoints in telephone exchanges. Leiserson (1985) at MIT generalized the concept as the "fat-tree," where the tree is "fat" because link bandwidth increases toward the root, proving it could emulate any network of equal hardware volume. The same cost-minimization logic that drove 1950s telephony now drives ML cluster design: minimize switch count while guaranteeing non-blocking connectivity for global AllReduce.

Napkin Math 3.2: Bisection bandwidth: The cost of oversubscription

Problem: A cluster designer is choosing between a "non-blocking" (1:1) fat-tree and a "cost-optimized" (4:1) spine for a $300M cluster of 1024 GPUs. This is a separate economics scenario rather than the running 175B job: each GPU contributes a 5 GB gradient payload, and the workload is assumed to spend $f = 30$ percent of each step communicating. That share is a

scenario input. It is larger than the 13 percent of the illustrative schedule in figure 3.2, and neither is a measured share of the reference job. How much slower will the AllReduce be on the cheaper network, and what does that cost?

Math: Bisection bandwidth ($\text{BW}_{\text{bisect}}$) is the minimum pipe diameter between halves of the cluster. Assume a unidirectional ring whose ranks alternate between the two halves of the cluster. Every ring edge then crosses the cut, with $N_{\text{GPU}}/2$ senders in each direction. Assume balanced routing, no competing traffic, and negligible startup latency. ReduceScatter plus AllGather contributes a traffic factor $2(N_{\text{GPU}}-1)/N_{\text{GPU}}$ = 1.998 per GPU; multiplying by $N_{\text{GPU}}/2$ gives the bytes crossing the cut in each direction. With a 5 GB payload, a traffic factor of 1.998, and 512 senders per direction, 5,115 GB crosses the cut. A ring that groups ranks within each half would have fewer crossing edges and a different cut-load bound.

1. **Non-blocking (1:1)**: $\text{BW}_{\text{bisect}} = 512 \times 50$ GB/s = 25,600 GB/s per direction.
   - Time: 5,115 GB / 25,600 GB/s $\approx$ 0.200 s.
2. **Oversubscribed (4:1)**: $\text{BW}_{\text{bisect}}$ = 25,600 GB/s divided by 4 = 6,400 GB/s per direction.
   - Time: 5,115 GB / 6,400 GB/s $\approx$ 0.799 s.
3. **Throughput**: If the communication share stretches by the slowdown factor r while computation stays fixed, then relative throughput $= \frac{1}{(1-f)+fr} = \frac{1}{0.70+0.30\times4} = \frac{1}{1.90} \approx 0.53$, a loss of about 47.4 percent of cluster throughput.
4. **Cost**: That loss times the $300M capital cost is about $142.1M.

Systems insight: For this cut-saturating placement, saving money on spine switches creates a 4× bottleneck for global synchronization, and on a $300M cluster that bottleneck idles about $142.1M of compute. The same formula gives the 23.1 percent loss of the 2:1 spine in the non-blocking definition. This is an equivalent-capacity cost under the stated workload assumptions, not a direct equipment loss. Oversubscription must be evaluated against the collective traffic actually crossing the constrained cut.

⇄ **Contrast:** The GPT-3 synchronization example in section 1.5.5 assumes pure data parallelism and no overlap, so its exposed communication share far exceeds the share assumed here.

The subscription ratio at each tier therefore determines whether a fat-tree can sustain global collectives, and the tier count determines how far it can grow. Scaling beyond the 2,048-host two-tier limit requires a three-tier architecture with core switches, enabling the fabric to reach 65,536 hosts, but the added scale carries both cost and latency. Each of the k pods contains $k/2$ leaf switches and $k/2$ spine switches (k^2 switches across all pods), and their $k/2$ spines each send $k/2$ uplinks into $(k/2)^2 = k^2/4$ core switches. A non-blocking $k = 64$ three-tier tree therefore requires $5k^2/4 =$ 5,120 switches across the core, spine, and leaf layers. At $10,000 to $50,000 per switch, the switching layer alone represents $51M to $256M. Typical hop counts also rise, from two inter-switch links and three switch traversals for intra-pod traffic (leaf-spine-leaf) to four inter-switch links and five switch traversals for inter-pod traffic (leaf-spine-core-spine-leaf), adding serialization delay and switch processing time to the α latency term across thousands of synchronization steps. Rail-optimized topologies respond to both pressures by matching the physical cabling pattern to the collective communication pattern.

Checkpoint 3.2: Fat-tree topologies

These questions check whether hierarchical switch-fabric trade-offs are clear:

- ☐ **Radix and scale**: compute the maximum number of GPUs a radix-64 two-tier fat-tree can connect without core switches.
- ☐ **Non-blocking ratio**: explain why a non-blocking fabric requires a 1:1 subscription ratio at every tier of the network.
- ☐ **Tier vs. latency**: determine how α (startup latency) changes when moving from a two-tier to a three-tier fat-tree.

- ☐ **Aggregate vs. bisection**: distinguish aggregate bandwidth from BW_{bisect} and identify the metric that matters for AllReduce.

3.5.3 Rail-optimized topology

When tensor-parallel groups stay inside each node's high-bandwidth NVSwitch domain, data parallelism exchanges gradients only between GPUs that hold the same model shard in the same local slot position across nodes, yet a standard fat-tree wires all eight GPUs of a node into a shared ToR switch. A **rail-optimized topology** instead cables external network ports by local slot position (figure 3.10): GPU 0 in every node connects to the Rail 0 switch, GPU 1 connects to Rail 1, and so on across the cluster.

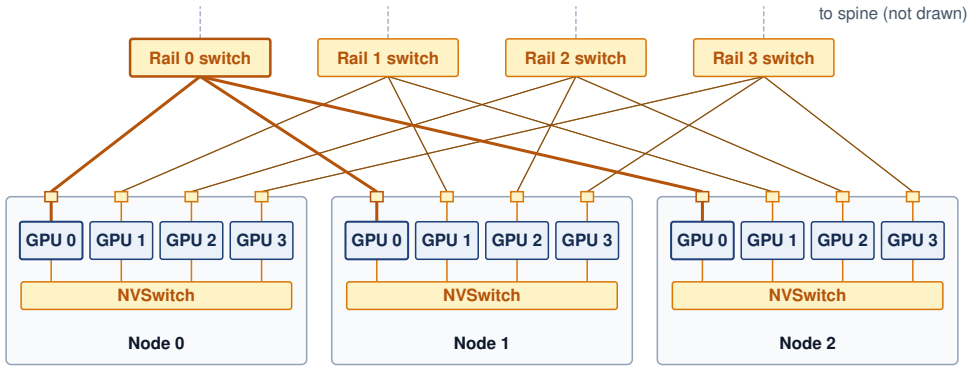

Figure 3.10: Rail-Optimized Physical Wiring: In dense accelerator nodes such as the DGX H100, each GPU's network port is cabled by local slot position: GPU 0 in every node connects to the rail 0 switch, GPU 1 to the rail 1 switch, and so on. Three nodes and four of each node's eight GPUs are drawn, with the rail 0 cables emphasized; further nodes attach to the same rail switches. Tensor-parallel traffic stays on each node's NVSwitch, same-rank data-parallel traffic stays on its own rail, and a spine layer (not drawn) bridges the rails for cross-rank traffic.

This slot-aligned wiring exploits the stratified traffic that 3D parallelism produces. Because each node holds one full tensor-parallel group, with one shard on each GPU, each data-parallel replica assigns the corresponding model shard to the same local GPU slot (where the slot number is a local rank within a node, distinct from the worker's global rank in the distributed job). Replicas are identical copies of one layout, so nodes that hold the same pipeline stage hold the same slice of the model in a given slot. During data-parallel AllReduce, GPU 0 on one node synchronizes its shard's gradients with GPU 0 on other nodes, but rarely communicates with GPU 1. Isolating each local rank onto its own rail switch fabric removes same-rank traffic from shared rack switches and, when a rail fits on a single switch, shortens the path between corresponding GPUs. The isolation holds however a rail is built, whereas the shorter path does not.

Napkin Math 3.3: The rail-optimized dividend

Problem: A team is synchronizing per-rank data-parallel gradients across 128 nodes. In a standard fat-tree, each message between same-rank GPUs traverses its source leaf switch, a spine switch, and the destination leaf switch (3 switch traversals). In a rail-optimized network built from rail switches with at least one port per node, all corresponding GPUs sit on the same rail switch (one traversal). On the reference cluster that requires rail switches with more ports than the radix-64 switch used so far, so the estimate describes a higher-radix design. How much latency dividend does the rail design earn when a rail fits on one switch?

Math: This estimate counts only the switch-traversal contribution to communication latency (α), using the same assumed delay for each traversal.

1. **Standard latency**: 3 traversals × 0.6 µs = 1.8 µs.
2. **Rail-optimized**: 1 traversal × 0.6 µs = 0.6 µs.
3. **Result**: 3× lower latency.

Systems insight: For synchronous data-parallel AllReduce, this 3× reduction applies to the modeled switch-latency contribution, not to total collective time. Its effect on scaling efficiency (η_{scaling}) depends on payload size, delivered bandwidth, compute time, and overlap. The single-switch rail in this example avoids the spine traversal for same-rank traffic. The dividend is conditional. On the reference cluster's radix-64 switches each rail outgrows one switch and regains the extra traversals, so what rails reliably buy there is isolation, developed below.

The hop saving is only part of the case for rails. Large language model fleets adopt the pattern because it matches how their parallelism plan divides traffic.

Lighthouse 3.1: Archetype A (GPT-4/Llama-3): The rail-optimized fleet

Archetype A's use of 3D parallelism generates two distinct traffic patterns that pull in opposite directions: (1) large, bandwidth-hungry gradient averaging for data parallelism, and (2) high-frequency, latency-sensitive activation exchanges for tensor parallelism. Keeping tensor parallelism inside each node's NVLink domain leaves data-parallel gradient traffic on the fabric, and rails give each same-rank data-parallel group its own switch network. A rail can provide a single switch hop between nodes when its endpoint count fits the available switch ports. Larger rails require additional switching stages, so rail alignment alone does not guarantee single-hop communication. Either way, the lesson for Archetype A is that how the fabric is wired to the parallelism plan matters as much as how many accelerators it connects.

For the cluster of 1,024 GPUs spanning 128 nodes, this creates eight parallel networks of 128 GPUs each, requiring at least one switch port per endpoint in each rail. A single-switch realization therefore needs a radix at least as large as the rail endpoint count; the earlier radix-64 switch cannot realize this rail alone, which is why notebook 3.3 assumed higher-radix rail switches. Built from radix-64 switches instead, each rail becomes a small two-tier fabric that preserves rank alignment but adds hops. For the bandwidth-hungry gradient exchanges of data parallelism, the larger benefit is isolation: each rail carries only its own rank's traffic, so the eight rail networks operate concurrently without competing for shared spine capacity. When a rail fits on one switch, the shorter path also trims the startup cost a ring pays on every round.

That same isolation creates a sharp trade-off, because traffic that must reach GPUs of different ranks requires bridging across rails. Such traffic includes expert routing in a mixture of experts (MoE) model, which sends each token to the few expert subnetworks chosen for it, and pipeline-stage handoffs that do not align with the rail wiring. Large clusters therefore often employ a hybrid approach, using rail-optimized leaves for same-rank data-parallel traffic while bridging the rails with a full fat-tree spine to support cross-rank communication patterns.

↪ Expert parallelism, in which each rank hosts a different subset of experts and tokens travel to them by AllToAll, is developed in section 6.5.5.

Checkpoint 3.3: Rail-optimized networks

These questions check whether workload-specific network-design trade-offs are clear:

- ☐ **Rail beneficiary**: identify the dimension of 3D parallelism (tensor parallelism, pipeline parallelism, or data parallelism) that benefits most from a rail-optimized design.
- ☐ **Switch hops**: explain when a rail-optimized network reduces the number of switch hops for gradient exchanges compared to a standard fat-tree, and what it still provides when a rail needs more ports than one switch offers.
- ☐ **Traffic isolation**: explain the traffic isolation benefit of placing all GPU 0s on one rail and all GPU 1s on another.

☐ **Core switch need**: classify the claim as true or false: a rail-optimized network eliminates the need for high-bandwidth core switches connecting different rails.

3.5.4 Dragonfly and torus alternatives

At tens of thousands of endpoints, a non-blocking three-tier fat-tree forces a third layer of core switches and tens of thousands of long optical cables into the budget, driving network capital expenditure into the tens or hundreds of millions of dollars. Two alternative topologies attack that cost wall from opposite directions: tapering the global switch tier or removing external fabric switches altogether.

A **dragonfly**[7] topology organizes high-radix routers (the dragonfly literature's name for switches) into fully connected groups joined by global optical links. Each group has a routers, and each router has p terminal ports and h global channels. The ratio of h to p sets how much global bandwidth each group receives relative to the traffic its terminals can inject, so global bandwidth is a design parameter rather than a fixed percentage of aggregate injection bandwidth. The balanced configuration analyzed by Kim et al. (2008), $a = 2p = 2h$, is one setting of that knob. A tapered dragonfly gives each router fewer global channels than terminal ports, trading cut capacity for fewer long optical links. The routers in a group are fully connected, and the groups are joined by global links. A minimal route between two groups therefore takes at most one local hop to the router holding the right global link, one global hop, and one local hop in the destination group. No minimal path exceeds three router-to-router hops, regardless of system size. Short paths come cheaply, but performance for jobs spanning groups depends on the chosen p/h ratio, routing, and traffic matrix, because every inter-group transfer shares the provisioned global links.

A **torus** topology connects each node directly to its neighbors in a multidimensional grid, most commonly a 3D torus where every node links to its six adjacent peers (up/down, left/right, front/back). The design offers full local bandwidth with minimal switching hardware, because every neighbor is a single hop away. However, global communication requires traversing the diameter of the mesh ($\mathcal{O}(N^{1/3})$ hops), making latency scale poorly with cluster size. Google adopted this architecture for its Tensor Processing Unit (TPU) pods because the collectives that dominate its training workloads can be scheduled as rings along each torus dimension, a nearest-neighbor pattern that maps naturally onto the physical grid. The limitation becomes apparent with MoE, which relies on AllToAll communication patterns. On a torus, this traffic can congest the limited $\text{BW}_{\text{bisect}}$ of the mesh and degrade performance relative to a non-blocking switch fabric.

The trade-off becomes concrete for the 4,096-endpoint deployment that the three-tier fat-tree reached with four pods. Arranged as a $16 \times 16 \times 16$ torus with wraparound links, the same endpoints need no external switching stages. Splitting the cube in half across one axis cuts a 16×16 plane of interior links and an equal plane of wraparound links, so only $2 \times 16^2 = 512$ links cross the bisection, against 2,048 in the non-blocking fat-tree. With 50 GB/s per link and direction, the torus offers 25.6 TB/s across the cut against the fat-tree's 102.4 TB/s, a 4× gap. Because wraparound links cap the distance along each 16-node axis at 8 hops, the torus's worst-case path across all three axes crosses $3 \times 8 = 24$ neighbor hops against the fat-tree's five switch traversals. Because a cubic torus's cut grows only as $N^{2/3}$ while its endpoint count grows as N, that capacity gap widens at fixed link rate.

Whether the gap matters depends on the traffic. A ring AllReduce scheduled along each torus dimension moves data only between neighbors and combines contributions before forwarding them, so it leans little on the narrow cut. An MoE AllToAll or an embedding exchange among a recommendation model's table shards instead sends a large share of its traffic between distant endpoints and drives it straight into the cut, with the exact pressure set by placement and message sizes. Consequently, the collective schedule and traffic matrix determine whether the torus's lower cross-sectional capacity limits completion time.

At the scale of the 1,024-GPU reference cluster, the four families range in cut capacity from 25.6 TB/s in the non-blocking fat-tree down to 6.4 TB/s in a 3D torus, while an assumed interconnect cost per Gb/s of bisection bandwidth runs from $1.5/Gb/s for the switchless torus to $4/Gb/s for the tapered dragonfly (figure 3.11). The cheapest capacity is also the least capacity, which is the trade every topology decision prices. The dragonfly's assumed cost per Gb/s is the highest here because its savings come from replacing long optical runs and switch tiers at large scale, which Kim

7 **Dragonfly topology**: Introduced by Kim et al. (2008) at ISCA 2008, the dragonfly uses high-radix routers grouped into fully connected "super-nodes." For systems with at least 16,000 terminals, the architecture achieves a 52 percent network-cost reduction relative to a folded Clos (the fat-tree of section 3.5.2) of equivalent bandwidth, though global optical link oversubscription can create AllToAll bottlenecks during MoE token routing.

↪ AllToAll schedules and their cost at scale, where every rank sends a distinct share to every other rank, are developed in section 5.8.

et al. (2008) reports for systems of at least 16,000 terminals; at 1,024 GPUs it has few such runs to replace. The rail's 12.8 TB/s applies to cross-rail traffic only, because same-rank traffic stays on its own rail. Radix, endpoint count, link speed, routing mechanisms, and equipment pricing all shift these trade-offs; the topology family alone dictates neither raw cross-sectional capacity nor total cost.

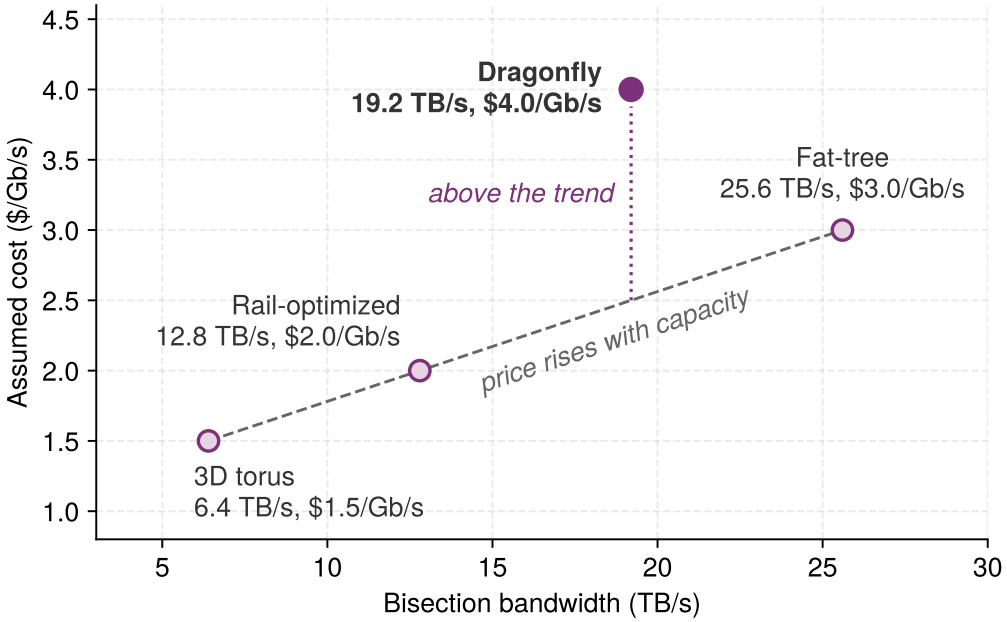

Figure 3.11: An Illustrative Bandwidth-Cost Scenario for ML Network Topologies: BW_{bisect} (TB/s, horizontal axis) against assumed cost per Gb/s (vertical axis) for four topology configurations on a 1,024-GPU cluster: fat-tree (25.6 TB/s, $3/Gb/s), dragonfly with tapered global links (19.2 TB/s, $4), rail-optimized with a 2:1 bridging spine (12.8 TB/s, $2), and a 3D torus (6.4 TB/s, $1.5). The bandwidths follow from the stated structures at one link rate; the costs are scenario inputs that illustrate the comparison method rather than universal properties of each topology family.

3.5.5 Optical circuit switching and reconfigurable topology

Every topology compared in figure 3.11 is fixed when the facility is cabled, locking the cluster into a single physical graph regardless of which workload runs on it. A statically cabled Clos hierarchy (section 3.5.2) fixes the oversubscription ratio and bisection pattern whether the active job is pipeline-dominated, data-parallel, or AllToAll expert-routed. A statically cabled torus (section 3.5.4) is equally rigid: it fits dimension-ordered rings, but it congests under MoE AllToAll traffic and cannot rewire its links to match the job.

Switching those multi-tier fabrics electrically also adds a recurring power cost at every hop. Each electrical packet switch inspects headers and chooses an outgoing link only after an Optical-Electrical-Optical (OEO) conversion: incoming photons hit a photodiode, high-speed SerDes circuits deserialize the electrical signal, switch ASICs buffer and route packets across crossbars, and output SerDes drive lasers to re-transmit light over optical fiber. At 800 Gb/s and 1.6 Tb/s per-port data rates, SerDes and transceiver electronics consume significant energy (section 3.3.3), and every additional switch tier repeats that conversion cost for each bit that crosses it. Changing the physical connections between node groups in light, for the duration of a workload allocation rather than packet by packet, avoids both static wiring and per-hop OEO conversion.

Definition 3.7: Optical circuit switching

Optical circuit switching (OCS) is a physical networking architecture that routes optical signals directly using Micro-Electro-Mechanical Systems (MEMS) mirror arrays without Optical-Electrical-Optical (OEO) conversion, enabling dynamic, software-defined topology reconfiguration and high energy efficiency at warehouse scale.

1. **Significance**: An optical circuit bypasses packet-by-packet electronic processing in the switching path. The complete system still consumes power for actuators, control, and transceivers. Its compatibility across line rates depends on optical loss, wavelengths, and endpoint budgets, so protocol transparency does not imply unrestricted upgrades.
2. **Distinction**: Unlike electrical packet switches (which inspect headers, buffer packets in SRAM, and make nanosecond routing decisions per packet), an OCS establishes semi-static physical optical circuits with millisecond-scale switching latency (10–30 ms), making it suitable for topology-level reconfiguration rather than dynamic packet routing.
3. **Common pitfall**: Assuming OCS replaces packet switching entirely. OCS operates at circuit granularity: it configures cluster-scale topologies (such as 3D tori or job-specific sub-clusters) and routes around hardware faults, while packet switching or accelerator-native interconnects handle fine-grained, intra-step message routing.

Google's TPU v4 supercomputer (Jouppi et al. 2023) deployed optical circuit switching in production. Instead of wiring 4,096 chips into an inflexible electrical fat-tree, the system groups accelerators into $4\times4\times4$ node cubes and routes the inter-cube links through an array of 3D MEMS optical circuit switches (codenamed Palomar).

The optical layer serves two purposes, adapting the topology to the workload and recovering from faults. For workload adaptation, the physical interconnect can be reconfigured to match the communication structure of the running model. A training run using 1D pipeline parallelism, whose stages form a line, can stretch its slice along one long axis to streamline stage-to-stage transfers. Conversely, a 2D or 3D tensor-parallel partition, which splits each tensor across a grid of chips, can configure a compact cubic 3D torus that minimizes average hop count, while an MoE workload can maximize bisection bandwidth across specific expert groups.

Optical switching also supports fault isolation by changing connections around unavailable hardware where spare connectivity permits. The capability matters because in a statically cabled torus, a single failed link or node can break a ring along an entire dimension and degrade every job that spans it. With optical circuit switching, the mirrors can be redirected to bypass the failed cube or fiber and splice healthy cubes into a complete, smaller torus slice without manual recabling. The fabric manager must first detect the fault, choose a feasible allocation, and establish the new circuits, which means circuit-switching time is only one part of recovery; a training job may also need to rebuild its process groups and restore state before computation resumes, so optical reconfiguration latency alone does not specify recovery time.

The trade-off is switching latency. Because MEMS mirrors physically rotate to steer collimated light beams, reconfiguration occurs on a mechanical timescale that is much longer than packet forwarding in an established electrical path. Circuit setup changes connectivity, while packet forwarding uses connectivity already in place. Consequently, OCS raises the power and protocol ceiling of the switching layer while leaving reconfiguration latency as the limiting constraint, so it is deployed for job-level topology allocation and fault isolation, not for per-message dynamic routing. Within that limit, OCS turns the choice of topology from a decision fixed when the fabric is cabled into one made for each job.

Each topology comparison so far has considered one traffic pattern at a time. In practice, clusters serve multiple workloads with different communication patterns, and topology selection must balance their competing demands.

Checkpoint 3.4: Topology selection for your workload

The choice of network topology dictates the upper bound of training efficiency. Warm up by matching each single-workload pattern to its ideal topology, then design a fabric that must serve several at once.

Single-workload picks

- ☐ **AllReduce bandwidth**: select the topology that maximizes BW_{bisect} for a standard data-parallel job dominated by AllReduce.
- ☐ **MoE traffic**: select the topology that best sustains a mixture-of-experts model's AllToAll-dominant, bursty communication, in which any rank may send to any other.
- ☐ **Neighbor locality**: select the topology that best fits a ring AllReduce scheduled along each dimension of the machine, where every transfer goes to a physical neighbor.

Designing for a mixed cluster

You are designing the network for a new ML cluster that will run two primary workloads: (1) training a 175B-parameter language model using 3D parallelism (tensor, pipeline, and data parallelism), and (2) serving a mixture-of-experts model that relies heavily on AllToAll communication to route tokens to the correct experts.

- ☐ **Training topology**: select the topology (fat-tree, rail-optimized, or hybrid) that best matches tensor parallelism within groups of 8 GPUs and data parallelism across all nodes, then justify the choice.
- ☐ **MoE topology**: revise the topology choice for unpredictable AllToAll traffic where any GPU may need to communicate with any other GPU.
- ☐ **Oversubscription risk**: identify which workload suffers more under 2:1 spine-layer oversubscription and propose a scheduling strategy to mitigate the impact.

Topology provides the structural capacity for the cluster of 1,024 GPUs to communicate, but structure alone does not guarantee performance. When the 1,024 GPUs synchronize contributions to the model's 350 GB gradient state, that theoretical capacity collides with the reality of **fabric behavior**: congestion control and routing dynamics determine whether this traffic flows smoothly or gridlocks under the strict synchronization of the BSP model.

3.6 Level 4 Fabric Behavior, Congestion, and Routing

A collective can finish late even when the topology provides full BW_{bisect}. Levels 1 through 3 determine how long a transfer takes on an unloaded path; fabric behavior determines how much longer it takes once the job's own synchronized traffic loads that path. Competing flows can concentrate on one output path, and the BSP barrier introduced in section 3.2 exposes any delay that extends the last required transfer. The relevant paths are those used by the participating workers; a congested link elsewhere in the fleet need not affect the job. In the Fleet Law (equation 1.3), measured transfer durations contribute to $T_{comm}(N)$, while residual waiting at synchronization contributes to $T_{sync}(N)$, the point where the communication axis of C^3 hands its cost to the coordination axis. A timeline must assign each interval consistently so the same congestion delay is not counted twice.

The p99 tail latency introduced with the BSP barrier helps expose delays that an average hides, but a particular barrier waits for its last required participant, not for a percentile. Evaluating a fabric therefore requires both transfer-time distributions and the collective completion times they produce under load.

3.6.1 Priority flow control (PFC)

The BSP weakest-link property means that even a single dropped packet can trigger a retransmission timeout that stalls the global barrier for milliseconds, a delay that can rival the entire communication phase of a training step. PFC-based RoCEv2 fabrics prevent the drop in the first place by running PFC, the pause mechanism introduced with the RDMA stacks (section 3.4.2), in lossless mode, so that upstream senders stop before any buffer saturates (Guo et al. 2016; Gangidi et al. 2024).

Definition 3.8: Priority flow control

Priority flow control (PFC) is a link-layer mechanism used by RoCEv2 ML cluster fabrics to prevent switch buffer overflow. When a port's queue depth crosses a configured threshold, the switch sends PAUSE frames to an upstream sender, which throttles injection on a per-priority basis without dropping packets.

1. **Significance**: In PFC-based RoCEv2 deployments, PFC is the foundation for a lossless Ethernet service. A PFC PAUSE frame must reach the upstream sender before the buffer overflows, so each pause threshold must leave headroom for the data still in flight during roughly one link round trip.
2. **Distinction**: Unlike standard IEEE 802.3 flow control, which pauses all traffic classes on a link, PFC operates per priority class, allowing latency-sensitive control messages to continue flowing while only pausing the congested gradient data queue.
3. **Common pitfall**: A frequent misconception is that PFC solves congestion. It transforms packet loss into congestion spreading. PFC itself places no bound on how far PAUSE frames propagate, so a single faulty transceiver can pause traffic far from its own link unless a watchdog intervenes.

The danger of PFC lies in its cascading nature. When a switch port's buffer fills, it sends a PAUSE frame upstream, which can move congestion into the upstream switch and trigger another PAUSE frame. A slow receiver, routing hot spot, or true fan-in burst (many senders converging on one port) can therefore pause links unrelated to the original queue. This behavior, known as **congestion spreading**, is the primary operational risk of PFC-based lossless Ethernet, and the flows it stalls without ever touching the congested queue are called **victim flows**. Fan-in is one trigger, but it is not intrinsic to every collective. A bandwidth-optimal ring exchanges data point to point, although multiple ring flows can still collide at a shared egress or an oversubscribed cut and create a hot spot without a single aggregation receiver (section 3.6.5).

3.6.2 Lossless fabric hazards and watchdog recovery

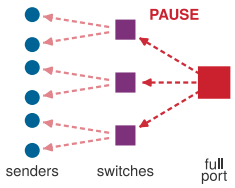

PAUSE frames climb back tier by tier and stall senders far from the full port.

While priority flow control provides the illusion of a lossless physical layer, relying on link-layer backpressure introduces structural failure modes that do not exist in loss-tolerant packet-switched networks. When switch queues cannot drop packets to relieve congestion, pressure propagates backward through the network graph, creating two distinct hyperscale failure modes: buffer-dependency deadlocks and PFC storms.

Definition 3.9: Cyclic buffer dependency deadlock

Cyclic buffer dependency deadlock is a permanent distributed lockup in a lossless network, whether its backpressure uses PFC PAUSE frames or InfiniBand credits. A set of switch port queues forms a closed dependency cycle, and each queue holds packets destined for the next queue while backpressure from that same queue pauses it.

1. **Significance**: A single such deadlock permanently halts traffic across all participating ports and upstream links, requiring manual intervention or watchdog timeouts that discard packets to recover. InfiniBand fabrics prevent it through deadlock-free routing and virtual-lane assignment, and PFC-based Ethernet fabrics need the same discipline in their routing and priority configuration.
2. **Distinction**: Unlike a PFC storm (which is a transient cascade of backpressure triggered by a slow-draining endpoint), a buffer-dependency deadlock is a static, self-sustaining topological cycle that cannot resolve itself even after traffic injection stops.
3. **Common pitfall**: Assuming multi-stage Clos (fat-tree) fabrics are inherently immune to deadlocks. Dynamic multipath routing, link failures, and address resolution flooding can introduce cyclic forwarding dependencies across switch tiers.

3.6.2.1 The graph mechanics of buffer-dependency deadlocks

In a lossless network, a switch cannot forward a packet to downstream buffer b_j if b_j is full; instead, it holds the packet in upstream buffer b_i and awaits credit or pause clearance. A directed buffer dependency graph $G = (V, E)$ formalizes this relationship. Each vertex $b_k \in V$ represents a packet buffer or virtual lane queue at a switch port. A directed edge $(b_i, b_j) \in E$ indicates that a packet currently occupying buffer b_i requires buffer b_j along its routing path to make forward progress.

A directed cycle in the buffer dependency graph is a necessary condition for a cyclic buffer dependency deadlock (definition 3.9), but a cycle alone does not establish that progress has stopped. The cyclic dependency in equation 3.1 becomes a deadlock when the occupied buffers cannot drain and every participant waits for its successor:

$$b_1 \to b_2 \to \ldots \to b_k \to b_1 \tag{3.1}$$

When each buffer in the cycle fills beyond its configured pause threshold, each port asserts a PFC PAUSE frame to its predecessor. Because no buffer can advance a packet into a paused successor, no buffer can drain, and no port can de-assert its PAUSE signal. The entire cycle freezes permanently. A four-switch ring whose queues (b_1 through b_4) are completely full shows this mutual protocol interlock (figure 3.12). Solid gray arrows trace packet forwarding dependencies advancing clockwise ($b_1 \to b_2 \to b_3 \to b_4 \to b_1$), while dashed red PAUSE frames propagate counterclockwise. Each full queue pauses its upstream predecessor while waiting for its downstream successor to drain, permanently freezing all four queues until an external watchdog intervenes to drop frames.

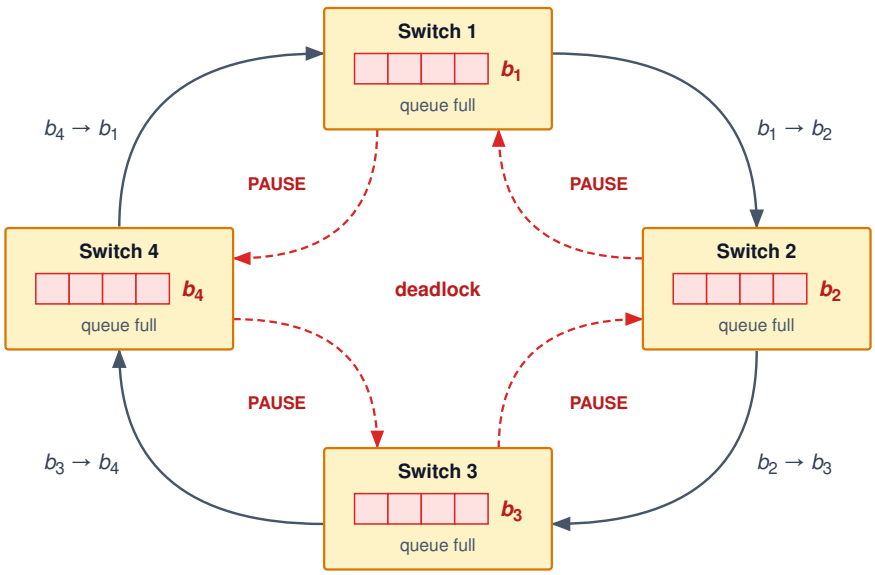

Figure 3.12: A PFC Deadlock from a Cyclic Buffer Dependency: Four switch queues form the cycle of equation 3.1. Follow the solid arrows clockwise: each full buffer holds packets that need space in the next. The dashed PAUSE frames run counterclockwise, so each full queue pauses its upstream predecessor while awaiting space in its downstream successor, and no queue can drain until a watchdog drops packets.

Routing changes and flooding can create the dependencies needed for deadlock. During link failure or route convergence, inconsistent forwarding tables can create transient loops. A lossy IP network eventually discards a looping packet when its time-to-live expires; a lossless class has no such release, so looping traffic can fill the buffers along the loop and hold them. If traffic occupies a cycle of mutually blocked queues, backpressure can prevent progress until recovery breaks the dependency. Multicast or broadcast replication can also fill queues whose dependencies intersect, while cross-connects and asymmetric failover paths can introduce cycles absent from the intended

hierarchical route. None of these conditions alone proves a deadlock; the occupied buffers and pause states must form the closed waiting cycle described by the model. Microsoft's production RoCE deployment encountered the flooding trigger directly (war story 3.1).

War Story 3.1: Microsoft's PFC deadlock (2016)

Context: A team led by Chuanxiong Guo at Microsoft Research deployed RoCEv2 across Microsoft's data centers, a fleet operating at the scale of tens of thousands of servers, to support latency-sensitive services. Their design provided a lossless fabric with PFC, and the team developed a Differentiated Services Code Point (DSCP)-based PFC mechanism to extend the deployment beyond virtual local area network (VLAN) boundaries (Guo et al. 2016).

Mechanism: PFC interacted with Ethernet flooding to create a cyclic buffer dependency. A dead server with an incomplete mapping between an IP address and a media-access-control address caused upstream switches to flood lossless packets to every port; PFC PAUSE frames then propagated upstream in response to the flooded traffic, forming a closed loop of paused buffers.

Impact: Lossless traffic through the switches in the cycle stopped entirely and stayed stopped, because no buffer in the dependency loop could drain without outside intervention.

Fix: Microsoft blocked broadcast and multicast traffic from entering lossless traffic classes, and dropped lossless packets when the corresponding Address Resolution Protocol (ARP) entry was incomplete rather than falling back to flooding.

Systems lesson: PFC backpressure is a structural vulnerability of lossless Ethernet at hyper-scale. RoCE fabrics require packet-class discipline, PFC pause telemetry, and hardware-level safeguards because a mechanism designed to avoid packet loss can itself become the failure mode. ML training fabrics built on RoCE inherit this risk directly. A single incomplete ARP entry can freeze the lossless class that carries a multi-thousand-GPU training run's gradients, with the cluster appearing healthy while no gradient flows between workers.

3.6.2.2 PFC storms and victim flows

In contrast to a static topological cycle, a **PFC storm** is a dynamic, runaway amplification phenomenon triggered when a single "slow drainer" (an accelerator stalled by memory bus saturation, operating system preemption, or a degraded transceiver with elevated bit error rates) fails to consume packets at wire speed. The attached ToR switch queue fills and sends a PAUSE frame to the sender. Because multi-port switches aggregate traffic from many sources across shared packet buffers, the paused queue blocks buffer allocation for unrelated packets. Packets of one priority class also leave the paused queue in order, so a packet bound for an uncongested port waits behind the blocked packets ahead of it (head-of-line blocking). The spine switches feeding that ToR then fill their buffers and pause the other leaf switches that send through them, and each additional tier widens the set of paused links until traffic between node pairs that never communicated with the stalled worker also stops. These victim flows turn one slow endpoint into a job-wide barrier delay, because every data-parallel worker whose path crosses a paused spine port now arrives late at the barrier.

Congestion spreading is the mechanism behind both a storm and a deadlock. A storm is the form that a slow drainer sustains, so it can end when the drainer recovers, unlike the closed cycle of a deadlock.

A degraded transceiver is one such slow drainer, so a component-exposure calculation shows how often a cluster should expect at least one, even though exposure alone does not estimate the probability of a PFC storm.

Napkin Math 3.4: Probability of at least one transceiver degradation

Problem: A 4096-GPU RoCE cluster, larger than the reference cluster, is in operation. If the probability that a given transceiver begins degrading on a given day is 0.001 percent, what is the chance that at least one transceiver degrades on a given day?

Math: Assume independent transceiver-degradation events with the same per-device probability.

1. **Total exposure**: 4096 × 3 link tiers = 12,288 links, or 24,576 transceivers.
2. **Probability of zero transceiver degradations**: $\Pr(\text{no degradation}) = (1 - 0.00001)^{24,576} \approx 0.782$.
3. **Probability of at least one degradation**: $\Pr(\text{at least one degradation}) = 1 - 0.782 = \mathbf{0.218}$ (about 21.8 percent).

Systems insight: Under these assumptions, a large RoCE cluster has a 21.8 percent daily chance of at least one degraded transceiver. This result motivates hardware telemetry and automated remediation, but estimating fabric-wide freeze risk additionally requires the conditional probability that a degradation triggers and propagates a PFC storm.

3.6.2.3 Watchdog recovery and the lossy alternative

Switches can use **PFC watchdog timers** to detect prolonged pause states and trigger recovery without a full reboot. The configured threshold T_{watchdog} must fit the deployment's pause and recovery budgets; a range such as 50–200 ms is a configuration example, not a universal requirement. Exceeding it indicates stalled progress, not proof of a cyclic deadlock. Depending on the implementation, mitigation can disable pause handling or discard affected traffic so blocked queues can drain. The transport must then recover any lost data, with the completion penalty determined by retransmission behavior, outstanding traffic, and the collective's dependencies. On an RNIC with Go-Back-N-style recovery (section 3.4.3), each discarded packet also costs the in-flight window behind it, so breaking a pause converts a stalled queue into a retransmission tail that the waiting collective absorbs.

This systemic brittleness is driving hyperscale operators toward next-generation architectures such as the Ultra Ethernet Consortium (UEC) (Ultra Ethernet Consortium 2025), detailed in section 3.10.1, and loss-tolerant RDMA designs (Mittal et al. 2018). Designs that combine packet spraying (sending one flow's packets over many paths) with congestion control and selective retransmission can operate without link-layer PFC. Removing PFC removes its pause-based deadlocks and storms, but endpoint recovery must still bound the completion-time penalty of congestion and loss.

3.6.3 End-to-end congestion control with DCQCN and HPCC

Because link-layer PFC fires only when a switch buffer is already near exhaustion, relying on pauses alone invites the cascades and watchdog drops of section 3.6.2. Keeping queues below those thresholds requires end-to-end congestion control, which throttles sender injection rates before buffers fill. For a BSP workload with a completion-time target, this rate control is a *latency requirement* rather than an *optional throughput optimization*: delaying a required transfer extends the synchronization boundary, so the controller must balance delivered throughput against queueing and recovery tails.

The first widely deployed solution, **DCQCN**,[8] operates as a reactive feedback loop using ECN (Zhu et al. 2015). When a switch's queue depth exceeds a configured threshold, it marks the ECN bit in the packet header. The receiver echoes this mark back to the sender via a congestion notification packet (CNP), prompting the sender to reduce its injection rate using a multiplicative-decrease algorithm. DCQCN is widely supported, but production AI collectives expose a tuning problem because the feedback loop takes a full network round trip to take effect. Setting the ECN threshold high leaves too little buffer headroom for synchronized bursts to slow down before hitting the PFC pause threshold, whereas setting it low causes senders to throttle aggressively on brief queue spikes and underutilize the link. Consequently, ECN thresholds and NIC firmware behavior trade off PFC avoidance, throughput, visibility, and tail latency rather than offering one stable optimum (Gangidi et al. 2024).

HPCC[9] uses in-network telemetry (INT) to give the sender more information about path load (Li et al. 2019). Switch measurements of queue occupancy and transmitted traffic let the controller estimate an appropriate injection rate rather than react only to a binary congestion mark. That information still arrives after a feedback delay, so the controller must account for changing load and avoid overreaction. The evaluated design reduces queue buildup and tail latency, but requires

[8] **DCQCN (data center quantized congestion notification)**: Introduced by Zhu et al. (2015) (Microsoft and Mellanox) at SIGCOMM 2015 for large-scale RDMA fabrics, its binary ECN/CNP feedback loop is simpler to deploy than telemetry-rich schemes. However, it creates a delicate tuning surface for training collectives: Meta's 2024 RoCE deployment showed that aggressive ECN marking reduces PFC pauses but degrades collective completion time (Gangidi et al. 2024).

[9] **HPCC (high precision congestion control)**: Introduced by Li et al. (2019) (Alibaba) at SIGCOMM 2019, HPCC requires programmable switch ASICs that append INT metadata to packets at line rate, a hardware dependency that limits the fabrics able to deploy it.

switch telemetry support and compatible endpoint processing. DCQCN needs less detailed switch feedback, trading that simpler deployment requirement against a different tuning problem.

3.6.4 Adaptive routing

Throttling senders with DCQCN or HPCC relieves a filling queue only by slowing the flows that cross it, even when parallel spine links sit idle. That imbalance is the ECMP hash-collision problem from section 3.2: because static flow hashing,[10] which works well for millions of small flows, pins each persistent elephant flow to a fixed path, an 8-flow AllReduce across 4 equal-cost paths can place 4 flows on one link while another carries none. In a synchronized training step, that overloaded link then determines completion time and can halve useful bandwidth.

Adaptive routing mitigates this problem by allowing switches to select among eligible output paths using congestion information rather than only a static hash. That choice moves traffic onto paths with room and adds no capacity, so it repairs a collision but cannot relieve a fabric whose paths are all full. The implementation is protocol- and device-dependent. Packet-level spraying[11] can send packets of one connection over different paths, so the receiver may observe reordering. Compatible endpoint support must accept out-of-order placement or restore the transport's required ordering; the presence of InfiniBand alone does not guarantee arbitrary spraying across every path. Ethernet fabrics can also employ **flowlet switching**, which changes a flow's path between bursts separated by a gap large enough to preserve packet order across the candidate paths, so its usefulness depends on how often the workload leaves such gaps relative to path delay differences.

MoE expert exchanges illustrate why path diversity matters: a dense exchange among 64 expert-parallel ranks contains at most $64(64-1) = 4{,}032$ directed remote rank pairs after excluding self-transfers, with active traffic governed by expert placement and token-routing sparsity. Where static-hash collisions concentrate that traffic onto a few links, adaptive routing spreads it across the provisioned $\text{BW}_{\text{bisect}}$. Archetype B (DLRM at Scale) creates similar exchanges when batches of embedding requests and results move among table shards, although batch and shard placement determine the traffic matrix rather than every individual lookup requiring a global exchange.

3.6.5 The incast problem in ML

Congestion control and adaptive routing manage flows that traverse the interior of the fabric. A different failure mode occurs at the edge, where the traffic pattern itself overwhelms a single port regardless of how much capacity the fabric interior provides.

Definition 3.10: Incast

Incast is a many-to-one ML cluster traffic pattern in which a large number of senders simultaneously transmit data to a single receiver port. It concentrates line-rate traffic from multiple sources into a single switch queue, and it can overflow the buffer if the burst outlasts the available buffering and control response, even when the rest of the fabric is uncongested.

1. **Significance**: Flat aggregation, synchronized storage reads, or other fan-in phases can direct many senders toward one receiver. The deliberately worst-case scenario here assumes a RoCE fabric whose ports all run at 400 Gb/s. In it, 256 senders each transmit at 50 GB/s toward one switch egress, producing 12.8 TB/s of offered load against a 400 Gb/s port. A 32 MB buffer then overflows in about 2.5 µs under the stated assumptions, after which the switch either drops packets or, in a PFC deployment, pauses its upstream neighbors. The example is a stress bound, not the traffic matrix of a ring AllReduce.
2. **Distinction**: Unlike general congestion (which occurs on shared internal links when aggregate traffic exceeds link capacity), incast is an endpoint bottleneck. It occurs even if every internal spine and leaf link is completely uncongested, because the bottleneck is the single destination port, not the fabric interior.
3. **Common pitfall**: Incast is not necessarily rare or unpredictable. Training jobs create it on a schedule, whenever many GPUs read dataset shards from or write checkpoints to the same storage node, and in the aggregation step of a tree or hierarchical reduction, where several ranks send partial sums to one parent. A gradient ring is different; its

[10] **ECMP (equal-cost multi-path)**: ECMP selects a path by hashing on the 5-tuple (source/destination IP, source/destination port, protocol), which means the same flow always takes the same path. This determinism is a feature for packet ordering but a liability for ML. Because a single AllReduce ring produces a small number of persistent flows, hash collisions are not transient statistical events but hot spots that last as long as those flows do.

[11] **Packet spraying**: Unlike standard Ethernet (which hashes flows to single paths to avoid reordering), packet spraying sends individual packets of a single flow across multiple paths. This can improve use of the available $\text{BW}_{\text{bisect}}$ by AllReduce elephant flows when the transport can tolerate or repair reordering. The Ultra Ethernet Consortium (UEC) specification design explicitly includes multi-path packet spraying and flexible ordering to reduce the ECMP collision problem for AI/HPC traffic (Ultra Ethernet Consortium 2025).

> point-to-point flows can still collide at a shared egress, but that is a hot spot rather than many-to-one incast. The severity depends on the collective algorithm, topology, routing, and oversubscription, so it must be measured and mitigated rather than treated as an edge case.

In a worst-case incast, 12.8 TB/s of offered traffic from 256 senders converges on a single 400 Gb/s (50 GB/s) egress (figure 3.13). Each sender injects at 50 GB/s. The buffer fills over a finite interval; control feedback can reduce arrivals, drops can discard excess traffic, or backpressure can move the queueing upstream. Whether that pressure spreads depends on shared queues and the paths carrying the traffic. Let the aggregate arrival rate to one egress be λ_{arr}, the egress service rate be μ, and the instantaneous queue occupancy be $Q(t)$. While a burst persists and $\lambda_{\text{arr}} > \mu$, the queue grows at the net rate

$$Q(t) = \min(Q_{\max},\, Q(0) + (\lambda_{\text{arr}} - \mu)t)\,,$$

where $Q_{\max}$ is the finite buffer capacity. The ratio $\rho_{\text{egress}} = \lambda_{\text{arr}}/\mu$ distinguishes a drainable workload from sustained overload: for $\rho_{\text{egress}} < 1$, idle intervals can drain the queue; as long-run ρ_{egress} approaches one, less slack remains to absorb bursts; and for $\rho_{\text{egress}} > 1$, occupancy must reach a control threshold or the buffer limit if the overload persists. This is a packet queue, distinct from the request queues of an inference server. Serialization and propagation set the unloaded path latency, whereas queued bytes add approximately Q/μ of waiting time.

⇄ **Contrast:** Request queues, whose arrival and service rates set serving latency rather than switch buffer occupancy, are developed in section F.1.1.

In the deliberately worst-case incast above, ρ_{egress} is 256×, so the 32 MB buffer fills in about 2.5 µs. Once full, that buffer represents 640 µs of egress waiting time. By contrast, the bandwidth-delay product (link rate times round-trip time) for the 400 Gb/s egress and an illustrative 5 µs feedback round trip is only 250 KB; it measures data in flight on the path, not the buffering needed to absorb an arbitrary fan-in burst. ECN must mark early enough for feedback to arrive before exhaustion, and in this worst case the buffer fills before a 5 µs round trip completes, so even a mark on the first queued byte returns too late. PFC can prevent a drop, but its backpressure can spread congestion. A larger buffer extends the reaction window at the cost of a larger tail-latency bound.

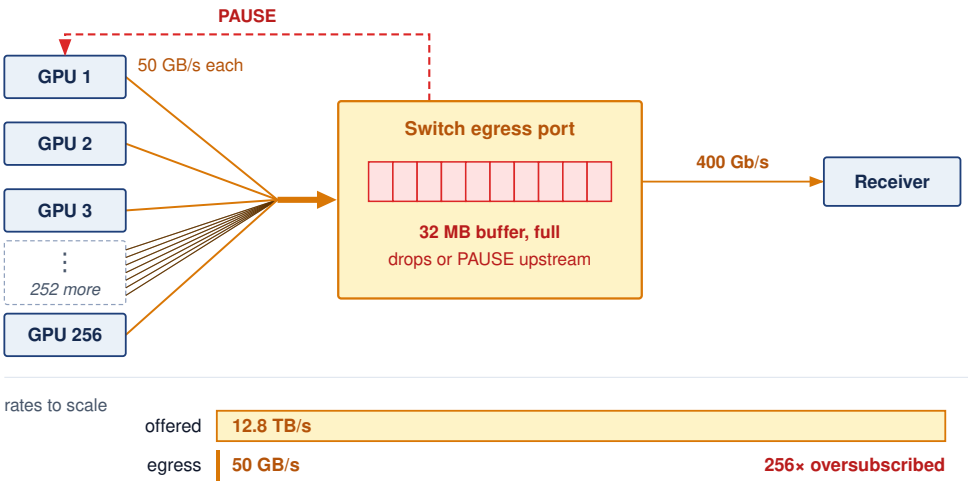

Figure 3.13: A True Incast Burst: This deliberately worst-case fan-in scenario sends 256 sources at 50 GB/s toward one 400 Gb/s (50 GB/s) egress with a 32 MB buffer. The 12.8 TB/s offered load oversubscribes the destination egress by 256×. Flat aggregation or synchronized reads can create this pattern; bandwidth-optimal ring AllReduce does not, because each rank exchanges point-to-point chunks.

Synchronized ML traffic is susceptible to both true fan-in and shared-egress hot spots. When a layer finishes backward computation, many ranks can inject traffic at once, and flat or hierarchical

aggregation phases can create true fan-in. Production clusters mitigate both cases along three complementary dimensions:

1. **Time.** Layer-staggering triggers each gradient bucket's AllReduce as soon as backpropagation produces it. Final-layer gradients reduce while earlier layers are still computing, so communication is spread across the backward pass instead of released as one synchronized burst. The same staggering is communication-computation overlap, so it also hides transfer time behind the backward computation that remains.
2. **Space.** Collective algorithm selection bounds fan-in at individual switch ports. Hierarchical and rail-local reductions aggregate gradients in stages so that no single port must absorb all senders at once, and topology-aware placement keeps concurrent ring flows off oversubscribed ports.
3. **Priority.** Quality of service (QoS) classification tags gradient traffic with the highest service class to reduce interference from background storage or management traffic; shared links, buffers, and arbitration still constrain isolation.

These mitigations spread the burst in time, space, and priority, but they do not eliminate congestion or tail latency. A production cluster must therefore deliver predictable step times with that residual congestion in place.

3.7 Level 5 Cluster Design and Case Studies

Cluster design reduces to concrete engineering decisions: which cables to buy, how to wire the racks, which protocol to deploy, and how to validate that the resulting fabric delivers the bandwidth the training job expects. At this final level of the stack, wires, transport, topology, and congestion control combine into a coherent system, an end-to-end gradient bus that makes thousands of distributed GPUs feel like a single machine. Two representative large-scale architectures, one built on InfiniBand and one on Ethernet, illuminate the trade-offs that define ML infrastructure (NVIDIA 2023; Meta Engineering 2024; Gangidi et al. 2024).

3.7.1 The GPU-to-GPU "gradient bus"

In a well-designed cluster, the network fabric acts as a scheduler-aware extension of the system bus, not a passive pipe. The ~9× per-direction cliff between NVLink and one external port, stated at the start of this chapter, is the gap the cluster must hide. Communication-computation overlap, which the Fleet Law credits as T_{overlap}, is the first mechanism for hiding it, through the same layer-staggered AllReduce launch that spreads incast bursts in section 3.6.5. Overlap cannot hide the last mile of the exchange, because the gradients computed at the end of backpropagation have little computation left behind which to hide their transfer (section 5.10).

Placement is the second mechanism, because it decides which communication crosses each boundary. The placement that the bandwidth hierarchy of section 2.5.1 sets, tensor parallelism inside the NVLink domain with pipeline stages and data-parallel replicas spanning nodes, becomes a wiring rule at cluster scale, because each fabric tier must be provisioned for the traffic placed on it. The links between nodes must meet the pipeline's activation-transfer budget, which physical adjacency alone does not guarantee. The full fabric carries the data-parallel exchanges, which gradient accumulation makes less frequent and overlap partly hides. Production designs therefore differ mainly in how they provision that outermost tier, as the two that follow show.

3.7.2 NVIDIA DGX SuperPOD architecture

The NVIDIA DGX SuperPOD, a concrete implementation of the gradient bus, provisions the outermost tier by obtaining losslessness from the link layer and aligning the wiring with the data-parallel ranks, so that the fabric's tail behavior depends on hardware mechanisms rather than on congestion tuning. It connects DGX H100 nodes using an NDR InfiniBand network (NVIDIA 2023). Each node acts as a dense compute island, with eight H100 GPUs connected via NVSwitch, which gives each GPU 900 GB/s of bidirectional internal bandwidth. Externally, each GPU pairs with a ConnectX-7 NIC delivering 400 Gb/s of injection bandwidth. Across a standard scalable unit of 32 nodes (256 GPUs), this yields an aggregate injection bandwidth of 12.8 TB/s (256 $\times$ 50 GB/s), the rate the scalable unit can offer the fabric when every GPU transmits at once.

This architecture connects all five levels of the network model into a single chain. At the wire (Level 1), it standardizes on 400 Gb/s NDR links, choosing copper or optics by the reach between each node and its leaf and between leaf and spine. Because those links run InfiniBand's native link-layer credit-based flow control (Level 2), they provide a lossless medium without the pause-cascade fragility of Ethernet PFC. Above the transport, the topology (Level 3) arranges the switches as a rail-optimized fat-tree, aligning GPU rails across each scalable unit so that same-rail traffic stays close while cross-rail traffic traverses the spine layer (replacing the simpler per-rack ToR allocation of section 3.5.1, which illustrated switch radix). When cross-rank traffic climbs to the spine, hardware-based adaptive routing (Level 4) selects among spine paths to preserve bisection behavior. Finally, at the cluster level (Level 5), additional switching joins multiple SuperPOD scalable units into larger fleets (NVIDIA 2023): for the 175B-parameter model, four SuperPOD scalable units wire together so that 1,024 GPUs function as one synchronous training system.

3.7.3 Meta Grand Teton platform

Meta's Grand Teton clusters take the opposite approach, accepting Ethernet's congestion-tuning burden in exchange for multi-vendor supply and shared operations. Meta disclosed two Grand Teton-based 24,576-H100 clusters: one using RoCE over an Arista/OCP Ethernet fabric and one using NVIDIA Quantum-2 InfiniBand, both with 400 Gb/s endpoints. Meta used both for Llama 3 training and reported ongoing Llama 3 training on the RoCE cluster (Meta Engineering 2024). Ethernet lets Meta source switches from multiple vendors and reuse the optical infrastructure and management tooling of its front-end serving fleet, avoiding the operational silo of a dedicated InfiniBand island.

Meta's RoCE study shows how Level 4 choices affect the complete fabric. It describes PFC watchdogs for long-duration pause events, iterative routing designs beyond plain ECMP, future-oriented flowlet switching experiments, and a 400G deployment in which DCQCN tuning proved difficult enough that Meta proceeded without DCQCN while relying on PFC plus higher-layer collective controls (Gangidi et al. 2024). That choice moved injection control from the NIC into the collective library rather than removing it, and the study cautions that the approach may not apply to all scenarios. The lesson is to validate the combination of endpoint behavior, routing, and flow control under the intended collective traffic. Large-transfer throughput alone does not establish small-message latency or recovery tails, and those results cannot be inferred from the Ethernet or InfiniBand label. Small-message latency matters in training through frequent small exchanges, which expose startup cost even when a full step lasts seconds, and in serving to a degree set by model partitioning, request shape, and service-level objectives (SLOs).

These case studies close the five-level stack by showing how the physical wire, transport protocol, topology, and congestion controls become one production fabric. The remaining operational problem is not how to build a single fast cluster, but how to share that expensive substrate across teams and workloads without sacrificing the predictable performance that training demands.

3.8 Network Virtualization

When a single training job leaves racks or spine links idle between runs, operators carve the cluster among multiple tenants and jobs, only to discover that a neighbor's checkpoint burst or evaluation sweep can starve a synchronous AllReduce across shared NICs and switch buffers. Network virtualization must make sharing compatible with bandwidth allocations, latency targets, and access isolation. These requirements interact because tenants still share physical links and buffers. An operator might reserve 80 percent of the cluster's BW_{bisect} for training and allocate the remaining 20 percent to inference and preprocessing. Enforcing that policy requires admission control and compatible scheduling along the relevant paths; a virtual NIC or separate traffic class alone does not guarantee the end-to-end share.

3.8.1 Hardware NIC virtualization with SR-IOV

Routing multi-tenant traffic through a hypervisor's software virtual switch reintroduces the exact host-CPU overhead and microsecond jitter that RDMA was built to bypass. Cloud providers instead expose NIC resources to virtualized GPU instances through **single-root I/O virtualization (SR-IOV)**, a standard that allows a physical NIC to present multiple **virtual functions (VFs)**. Each VF

exposes its own hardware queues, doorbell registers (which software writes to tell the NIC that work is waiting), and DMA mappings, and software can assign each VF to a workload. A directly assigned VF can remove host software mediation from the active DMA path, while drivers and virtualization software still configure access and manage the device. Performance depends on the NIC, I/O placement, and competing traffic. Because VFs share physical resources, placement logic must reserve NIC capacity and QoS policy alongside GPUs rather than treat the network as an unbounded pool.

However, hardware-level isolation still needs explicit bandwidth policy. SR-IOV exposes multiple VFs that share the NIC's physical resources; administrators can add per-VF or per-group QoS policies to cap or guarantee bandwidth, such as assigning eight VFs 50 Gb/s each on a 400 Gb/s NIC. For the 175B-parameter model training on a multi-tenant cluster, that configured partitioning can prevent a neighboring tenant from stealing bandwidth, but it also enforces a hard ceiling on peak throughput. The training job must be architected to operate within this slice, as no amount of software optimization can burst beyond the configured VF limit.

3.8.2 Traffic isolation and quality of service

Per-VF rate caps bound what each tenant injects at the NIC, but traffic from different tenants still merges on the fabric's shared switch buffers and spine links. When the 175B model is midway through a gradient AllReduce and a neighboring job initiates a 100 GB checkpoint save, that burst occupies a 400 Gb/s port for 2 s at full line rate. It can also claim a large share of the shared spine links for that interval, which stretches every AllReduce whose traffic crosses them. To prevent this noisy-neighbor effect inside the fabric, high-performance switches rely on **quality of service (QoS)** mechanisms that enforce fairness at the packet and queue level.

InfiniBand isolates traffic through virtual lanes (VLs), which maintain separate credit and buffer state within a physical link across up to 16 virtual lanes (one of them reserved for subnet management) (InfiniBand Trade Association 2000). If a storage burst exhausts the credits of its own lane, a gradient lane with available credits can continue transmitting, even though both lanes still share link bandwidth and arbitration. Ethernet and RoCE fabrics achieve a similar separation through device-dependent **traffic class** mappings and queue policies. Within those mappings, **enhanced transmission selection** allocates a share of link bandwidth to each traffic class, and switches may also partition buffer and forwarding resources into **network slices**. Across both stacks, these mechanisms reduce interference only when classification, resource allocation, and scheduling agree along the entire path, because shared priorities or buffer pools can still couple tenants.

Virtualization solves the sharing problem but makes performance diagnosis harder. When a training job slows down on a multi-tenant cluster, the cause could be a physical link degradation, a noisy neighbor exceeding its bandwidth allocation, or a misconfigured QoS policy. Distinguishing these cases requires systematic monitoring.

3.9 Monitoring and Debugging

Network performance problems in ML clusters are insidious because they show up as **silent waste** rather than explicit failures. A degraded transceiver causing a 10 percent reduction in effective bandwidth might slow each training iteration by only 2–3 percent. The slowdown is small because the lost bandwidth stretches only the exposed communication share of the step (the 30 percent assumed in the oversubscription estimate, less whatever overlap hides), and stretches it by roughly the inverse of the bandwidth that remains. A drift that small is easily masked by the natural variance of checkpointing or data loading. Over a 30-day training run on 1,024 GPUs, this invisible drag accumulates to roughly 15,000–22,000 wasted GPU-hours, burning about $30,000–$44,000 at $2/GPU-hour without triggering a single alarm.

Traditional IT monitoring tools like SNMP or ICMP ping measure connectivity, not the sustained throughput required by RDMA. Effective observability requires a three-layer approach: **application monitoring** (NCCL algorithmic bandwidth), **transport monitoring** (PFC pause frames, retransmission rates), and **physical monitoring** (FEC errors, signal attenuation) (Jeaugey 2017; Gangidi et al. 2024). The three layers retrace the five-level model from the top down. Application monitoring observes the collective that Level 5 integrates, transport monitoring observes the Level 2 recovery and Level 4 congestion machinery, and physical monitoring observes the Level 1 wire. Only by

correlating signals across these layers can operators detect that a "slow training run" is caused by a single degraded cable in one rack.

3.9.1 Link-level telemetry

A marginal cable or overheating transceiver rarely drops the link outright; instead, it accumulates symbol errors, triggers transmit discards, or retrains to a slower link state while the job keeps running. Catching that physical drift requires polling switch and NIC port counters, which report transmitted and received data, transmit discards, symbol errors, and link-down events; on InfiniBand a utility such as `perfquery` reads them, and Ethernet/RoCE devices expose the equivalent telemetry through their own management interfaces (Linux man-pages project 2026). Data counters measure utilization, discards warrant investigation against the port's traffic and recovery policy, and symbol errors and link-down events are evidence of a degrading signal. Correlating those counters with link rate, lane state, and transport behavior narrows the cause; a counter alone does not identify a defective cable or transceiver.

Polling link state alongside traffic counters can reveal **silent degradation**. A link operating at HDR (200 Gb/s) rather than the expected NDR rate (400 Gb/s) can remain connected while offering half the expected bandwidth. A link that trains with two of its four lanes active after a connector fault fails the same way: it stays up while its bandwidth is halved. Both cases are visible only when negotiated rate and lane state are compared with the deployment baseline.

For 1,024 GPUs and roughly 3,000 links, the assumed 0.1 percent point-in-time degradation probability gives 3 degraded links in expectation and, under independent events, a 95 percent probability of at least one degradation. This probability measures something different from the daily figure in notebook 3.4. That notebook asks how often a new degradation begins; this estimate asks how many degraded links are present at any moment. Standing prevalence, the share of links degraded at any moment, is roughly the onset rate times how long a degraded link stays in service, and silent degradation, by definition, stays in service until someone looks, so prevalence can far exceed what a small daily onset rate suggests. Both rates are scenario assumptions rather than measured fleet data.

A job slows when an affected path limits its required traffic; fleet-wide exposure alone does not determine the job's slowdown. When the degraded link lies on a ring AllReduce's path, however, the cost reaches every participant, because each round waits for its slowest hop and the whole ring proceeds at the pace of that one link.

↪ Gray failure, the general case of a component that degrades without failing, and why a synchronous job feels it as a straggler, is developed in section 7.9.5.

3.9.2 Bandwidth and latency validation

Clean port counters and an active 50 GB/s NDR link state still do not guarantee that a node pair can deliver that rate to a training job: a degraded PCIe path, a cross-socket NUMA bounce, or a congested spine cut can throttle RDMA transfers without incrementing a single physical error counter. Active validation must therefore measure the bandwidth and latency the job can actually achieve across the complete host-to-host path. Pairwise tests such as `ib_write_bw` can populate an **all-pairs bandwidth matrix**, or a sampled subset for a large fleet (linux-rdma project 2026). A low-throughput pair identifies a path to investigate. Comparing the affected pairs with routing and physical telemetry helps distinguish a shared switch or cable problem from endpoint limitations.

For latency-sensitive synchronization, `ib_write_lat` runs a ping-pong test with small RDMA writes and reports half the measured round-trip time as one-way latency. A reading such as 2 µs, 5 µs, or 100 µs must be compared with a baseline for the same endpoints, message size, path, and load, because no single value is a universal fault threshold or proof that traffic has fallen back to a slower transport. Application-level tests such as `nccl-tests` then exercise collective communication across the intended worker group. Point-to-point results establish a path baseline, while collective tests reveal costs from concurrent traffic and scheduling that isolated pairs cannot expose.

3.9.3 Systematic debugging workflow

When a training job reports lower-than-expected throughput, the diagnostic order should preserve the layer model rather than starting with cables. The following sequence moves from application symptoms toward physical causes through the three monitoring layers, with steps 1 and 2 at the application layer, steps 3 and 4 at the transport layer, and step 5 at the physical layer:

1. **Check GPU utilization**: Rule out compute bottlenecks with accelerator telemetry (`nvidia-smi`, for example). High SM utilization alone does not rule out a network bottleneck, because communication kernels can also keep SMs active. Correlate utilization with the compute and collective portions of the step timeline. Low utilization does not automatically implicate the fabric, however. A starved input pipeline (data loading or storage) produces the same symptom, so confirm the data path is keeping the accelerators fed before investigating the network.
2. **Inspect collective-library logs**: Enable the library's debug output (`NCCL_DEBUG=INFO` in NCCL, for example) to reveal which network transport was selected, the detected bandwidth between nodes, and any fallbacks to slower protocols.
3. **Run point-to-point tests**: Run a pairwise bandwidth test (such as `ib_write_bw`) between specific nodes in the job. A single degraded link can bottleneck the entire ring in a ring AllReduce.
4. **Check PFC/ECN counters**: Inspect switch counters along the path. Sustained PFC activity indicates persistent congestion that should be investigated at the scheduler or routing level.
5. **Validate physical layer**: Check symbol errors and CRC counts to identify failing transceivers or cables.

In this order, application symptoms and collective logs show whether the network is at fault and whether the library fell back to a slower protocol, point-to-point tests identify the failing path, transport counters show whether congestion is involved, and physical telemetry confirms whether the wire itself is degrading. Each layer therefore narrows the search before any hardware is replaced.

Checkpoint 3.5: Diagnosing a training slowdown

Your 175B model training job has been running for 3 days on the 1,024-GPU reference cluster. You notice that the iteration time has gradually increased from 4.2 s to 4.8 s (a 14 percent slowdown). The GPU utilization reported by `nvidia-smi` has dropped from 92 percent to 85 percent.

- ☐ **First diagnostic**: choose the first diagnostic step, name the tool that confirms whether the bottleneck is communication-bound or compute-bound, and cite the evidence that distinguishes the two.
- ☐ **Physical layer**: identify physical-layer causes that could explain a 50 percent bandwidth drop when `ib_write_bw` finds one node pair at 24 GB/s instead of the expected 48 GB/s, and verify them using switch counters.
- ☐ **Collective impact**: estimate the impact on overall AllReduce time when the degraded link lies in the ring AllReduce path, then evaluate whether a hierarchical reduction that routes less of the collective through that link would help.

3.10 Network Technologies That Move the Ceiling

Monitoring keeps deployed fabrics honest, but it cannot push past the physical ceilings it measures, such as the reach and power limits that copper approaches at high signaling rates. New fabric technologies earn attention only when they move one of those bounds. As clusters scale toward 100,000 GPUs, the durable design variables are power per bit, reach, and the ratio of compute to memory to network capacity, not the product names attached to a particular roadmap.

3.10.1 Protocol convergence and link density

Standard RoCEv2 leaves operators with fragile PFC pause thresholds and with ECMP hashing that was never designed for synchronized elephant flows, while InfiniBand locks the fabric into a single-vendor stack. The Ultra Ethernet Consortium (UEC) addresses these reliability and load-balancing limits for AI/HPC traffic over Ethernet. Its Ultra Ethernet Transport design combines packet spraying across multiple paths, flexible ordering, multiple transport delivery services, congestion-control changes, and telemetry (Ultra Ethernet Consortium 2025). The architectural question is how the

selected service delivers the required completion and ordering semantics under congestion or loss. Packet trimming (cutting an overflowing packet down to its header instead of dropping it, so the receiver learns of the loss at once), retry, and endpoint recovery should therefore be judged by their effect on completion-time tails and resource cost. Whether the design reproduces a lossless-fabric mechanism is not the test.

Higher switch capacity attacks hop count as well as raw bandwidth. XDR InfiniBand pushes 800 Gb/s per port, and Ethernet roadmaps include 800GbE and 1.6TbE link speeds (NVIDIA 2026f; Ethernet Alliance 2025). Ethernet switch silicon with 1.6 Tb/s ports pushes aggregate capacity to 102.4 Tb/s in current Tomahawk 6-class designs (Broadcom Inc. 2025). A single 1.6 Tb/s port delivers 4 times the bandwidth of a 400G port, so the same ASIC can instead expose 4 times as many 400G ports. Spending that capacity on radix rather than port rate is what flattens the topology, because fat-tree host count grows as $k^2/2$ at two tiers and $k^3/4$ at three. A cluster that previously required three tiers of switches may then be served by two when radix, port rate, and placement constraints line up, potentially reducing the number of optical links and switching stages. The transceiver savings require explicit port and cable accounting; removing one tier does not automatically halve the module count, and the latency benefit depends on the resulting paths and traffic.

3.10.2 Optical interconnects

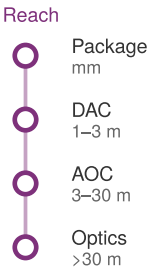

Each extra meter pushes the fabric from copper toward optics.

Pushing ports from 400 Gb/s to the 800 Gb/s and 1.6 Tb/s rates of section 3.10.1 takes more lanes, faster lanes, or both. Eight lanes in the 100 Gb/s class of current NDR links reach 800 Gb/s, but an 800 Gb/s port on four lanes or a 1.6 Tb/s port on eight lanes doubles the per-lane rate to the 200 Gb/s class, which sharply increases the insertion loss per meter of passive copper. As that usable reach shrinks, in line with the link-budget example in section 3.3.3, the SerDes must spend more energy per bit on equalization to recover the signal. Active electrical cables, which add retimer chips to regenerate the signal mid-cable, extend that reach but add latency and power of their own, leaving optical interconnects as the primary lever on power per bit and reach.

Co-packaged optics (CPO) places optical engines close to the switch or accelerator on a shared package substrate. Moving the optical engine from the switch faceplate to the package shortens the SerDes-to-optics electrical path from centimeters of board trace to millimeters, which reduces the signal-conditioning burden before transmission over fiber. The receiving optical engine converts the signal back to electrical form. Fiber still introduces propagation delay and has a loss budget; packaging the optics closer to the processor does not remove those constraints.

Co-packaged optics therefore changes the electrical portion of the link budget while retaining the reach advantages of optical transmission. Usable reach still depends on the fiber, connectors, optical power, and receiver design. The relevant comparison holds payload capacity and error targets fixed, then measures the power saved by the shorter electrical path against the packaging and maintenance costs of integration. The power side of that comparison can be estimated port by port for a single switch.

> Napkin Math 3.5: The optical dividend
>
> **Problem**: How much optics power does a 51.2 Tb/s switch save by moving from pluggable transceivers to co-packaged optics?
>
> **Math**:
>
> 1. **Pluggable architecture**: 128 ports $\times$ 20 W = 2.56 kW for optics alone.
> 2. **CPO architecture**: 128 engines $\times$ 10 W = 1.28 kW.
> 3. **Savings**: The savings reach 1.28 kW of power per switch.
>
> **Systems insight**: Across 1,000 switches, the assumed pluggable modules consume 2.56 MW. The modeled CPO configuration saves 1.28 MW, equivalent to the reference power budget of roughly 1,800 H100 GPUs. That equivalence measures the scale of the saving; using it for additional accelerators also requires compute capacity, power delivery, and cooling in the right locations. Optics power constrains both operating cost and the equipment a facility can support.

The spine example for 10,000 GPUs assigns 400 kW to pluggable optics. That subtotal counts only the selected spine-side endpoints. Applying the same fractional saving to those endpoints would require a CPO implementation with the assumed per-port power and capacity; opposite-end optics remain outside the comparison. Integration also changes serviceability: replacing a failed optical component may involve a larger assembly than swapping a pluggable module. The energy benefit must be evaluated alongside packaging, repair, and spare-capacity costs.

For ML infrastructure, CPO could flatten the bandwidth hierarchy by narrowing the gap between intra-node and inter-node bandwidth. If inter-node links achieve bandwidth comparable to NVLink-class local fabrics (hundreds of GB/s per GPU), the bandwidth penalty for extending tensor parallelism beyond a single node would shrink. Cross-node tensor parallelism would then become economical for more workloads, provided that startup latency and collective frequency also fit the step-time budget.

However, CPO introduces new challenges. Optical components are sensitive to temperature (laser wavelength shifts with temperature, requiring active thermal management), and integrating photonics on the same package as a 1,000 W GPU creates a hostile thermal environment for the optical components. The manufacturing processes for silicon photonics and CMOS transistors are similar but not identical, requiring separate fabrication steps that increase package cost and complexity. Hybrid integration, which places the photonics chiplet on a cooler region of the package substrate and thermally isolates it from the GPU die, is one way to manage that environment, at the cost of a longer electrical path than full co-location.

3.10.3 Disaggregated and composable architectures

As the gap between package-level and rack-level interconnects narrows, architects can reconsider the fixed compute-to-memory ratio of the server node, which leaves compute or memory underused when workloads need different resource ratios. Compute Express Link (CXL) memory pooling allows supported hosts to access assigned memory capacity through a CXL path, while disaggregated designs make some resource allocations independent of the chassis boundary. This changes which links lie on the workload's critical path. It does not merge CXL memory access, NVLink communication, and InfiniBand transport into one interchangeable service; each path retains its own supported operations, bandwidth, and latency.

A composed node therefore needs a new traffic budget. The α-β distinction between startup and transfer cost remains useful, but each path requires its own parameters and access model. Remote memory traffic can compete with other transfers, and fine-grained dependent accesses can expose latency that a bulk-bandwidth estimate hides. Disaggregation improves resource flexibility only when these added costs fit the workload's execution budget. Composition relocates the synchronization backbone onto new paths rather than eliminating it.

↩ CXL memory pooling, which makes node memory capacity less rigid but remains a capacity tier rather than a high-bandwidth memory (HBM) replacement, is established in section 2.9.1.

3.11 Fallacies and Pitfalls

Designing and operating high-performance fabrics for ML requires unlearning assumptions from traditional data-center networking. The following fallacies and pitfalls capture the most common errors that stall training and degrade cluster productivity.

Fallacy: *More bandwidth always means faster training.*

Engineers assume upgrading from HDR (200 Gb/s) to NDR (400 Gb/s) will yield proportional gains, but the α-β model (section 3.4.4) reveals this is only true in the bandwidth-dominated regime. For small models or the small-message phases of pipeline parallelism, the startup term α (1.5 µs in this chapter's NDR scenario, including NIC processing, switch hops, and FEC) dictates performance. If a workload is latency-bound, a 2× bandwidth increase may deliver only a small fraction of the line-rate gain while adding substantial power and transceiver cost.

Consider a 10 KB message (typical for control synchronization). On 200 Gb/s InfiniBand, it takes 1.90 µs. Upgrading to 400 Gb/s reduces this to 1.70 µs, a 10.5 percent improvement despite doubling the link rate.

Fallacy: *Lossless Ethernet is as reliable as InfiniBand.*

RoCE over Ethernet can achieve throughput comparable to InfiniBand for large transfers, but lossless operation depends on the configuration and interaction of PFC, ECN, routing, buffers, and NIC behavior (section 3.6.1). A slow drainer or firmware defect can trigger a PFC storm, where PAUSE frames spread beyond the original queue, and a forwarding loop or cyclic buffer dependency can freeze the lossless class outright. InfiniBand's per-hop credit mechanism avoids this specific reactive PAUSE-frame cascade, but the fabric still requires deadlock-free routing, sound virtual-lane configuration, and congestion management (InfiniBand Trade Association 2000). Both fabrics require validation and telemetry; RoCE adds particular sensitivity to the configuration of Ethernet losslessness. Procurement should therefore compare complete transport stacks rather than link rates, because the stack determines where flow-control state and operational risk reside even when nominal rates match.

Pitfall: *Accepting data-center oversubscription ratios for synchronous training.*

Oversubscription relies on the traffic crossing a shared boundary fitting its available capacity. Bulk-synchronous training can violate that assumption with coordinated bursts. In the cut-saturating scenario of section 3.5.2, 4:1 oversubscription increases the modeled communication bound by 4×. Applied to a $300M cluster with the assumed 30 percent communication fraction, the resulting idle-capacity estimate exceeds $142.1M. This is an equivalent-capacity comparison under the stated schedule, not an equipment loss or a universal penalty for every collective.

Fallacy: *TCP bandwidth tests are sufficient for ML fabric validation.*

Standard benchmarks like `iperf` measure kernel-based TCP/IP performance. TCP throughput depends on host processing, offloads, and tuning as well as wire speed; it has no fixed throughput ceiling independent of that configuration. An RDMA training job exercises a different data path, so a low TCP result does not establish that the RDMA path is slow, and a high TCP result does not validate collective performance. Validation protocols must use RDMA-specific tools from the `perftest` suite to match the workload's data path.

Pitfall: *Relying on adaptive routing instead of topology-aware placement.*

Adaptive routing can distribute traffic across available paths, but it cannot increase their total capacity. For a 1,024-GPU job spanning two oversubscribed spine groups, the group-to-group links become limiting when the collective sends more crossing traffic than those links can carry within its time budget. Placement can reduce that traffic, and routing can reduce path imbalance. Both must be evaluated against the actual collective schedule.

↪ Topology-aware gang placement that keeps collective traffic off contended cross-rack paths is developed in section 8.6.2.

Fallacy: *Network fabric problems always appear as hard failures.*

Network performance issues in ML clusters often show up as subtle training slowdowns rather than errors, because jobs continue to run while gradients arrive late. The silent-waste estimate in section 3.9 shows how a small bandwidth loss accumulates into wasted GPU-hours over a long run without raising a single alarm. Detection therefore depends on comparing delivered throughput with a baseline, not on waiting for an error.

Pitfall: *Treating PFC and ECN counters as optional production details.*

Without discard counters and PFC pause durations, queue pressure and recovery activity can stay hidden behind a slow collective, so both belong in production telemetry, correlated with collective completion time and the affected paths. Rising values do not by themselves establish a failing transceiver or predict an inevitable job failure. Physical error counters, traffic load, routing, and configured watchdog actions help distinguish the causes and select a response.

3.12 Summary

The five-level model of figure 3.3 bridges physical reach, transport primitives, switch topologies, congestion dynamics, and cluster integration into a single system contract. A fabric must provide both enough capacity for the workload's traffic and predictable completion at its synchronization boundaries. InfiniBand and RoCE distribute this responsibility differently across links, endpoints, and

operational policy; the production cases show why neither port rate nor protocol name determines delivered performance.

Three numerical anchors show how to apply the model. One GPU's NVLink bandwidth exceeds its own external port's per-direction bandwidth by 9×; the modeled NDR path has equal startup and transmission costs near 75 KB; and the 4:1 oversubscribed cut increases the fixed-payload transfer bound by 4×. Each result depends on its stated configuration and traffic assumptions. In the Fleet Law, transfer time contributes to $T_{\text{comm}}(N)$ and residual synchronization waiting to $T_{\text{sync}}(N)$; consistent timeline accounting prevents counting one delay in both. The communication and coordination axes of C^3 meet where those dependencies determine useful compute progress.

The same model guides design and diagnosis. Message sizes determine whether startup latency or bandwidth dominates; communication patterns determine which paths must carry that traffic. During operation, application throughput, transport counters, and physical telemetry distinguish inadequate capacity from contention or a degraded component. These measurements also test whether a proposed optical or composable architecture removes the limiting constraint or merely moves it.

Key Takeaways: The fabric decides useful compute

- **Synchronous training inverts data-center traffic**: A bulk-synchronous step releases a few large, correlated flows that must complete before the next step can start, so the fabric is judged by collective completion time and its p99 tail rather than by average per-flow throughput. Static ECMP hashing turns flow collisions into persistent hot spots.
- **Link speed is not fabric speed**: For traffic that crosses a fabric cut, available capacity bounds transfer time (principle 6). Workload placement and collective schedules determine how much data must cross, so advertised port speed alone cannot predict useful training throughput.
- **Latency and bandwidth limit different messages**: The α-β model of section 3.4.4 separates startup cost from per-byte transfer cost, with $n^* = \alpha \cdot \beta$ marking the regime change. Message-size distributions, not vendor peak numbers, determine whether to optimize software latency or hardware bandwidth.
- **Reliable delivery moves the risk**: RDMA transports need reliable completion and bounded recovery, not one universal lossless-fabric mechanism. InfiniBand normally suppresses loss with link-level credits; many RoCE deployments use PFC and ECN, while PFC-free designs place more recovery responsibility at the endpoints. Each choice trades hardware flexibility against tail latency and operational complexity.
- **Topology must match traffic**: Fat-trees buy flexible bisection bandwidth, rail-optimized designs give same-rank AllReduce traffic its own isolated switch network, and dragonfly or torus designs trade cabling and locality differently. The right fabric depends on whether the workload stresses AllReduce, AllToAll, or multi-tenant sharing, and optical circuit switching lets a pod revisit that choice per job at millisecond reconfiguration cost.
- **Losslessness and load create their own failures**: PFC prevents drops by pausing senders, which can spread one slow drainer's congestion to victim flows or freeze a lossless class in a cyclic buffer dependency. Under synchronized load, end-to-end congestion control is a latency requirement, adaptive routing repairs hash collisions without adding capacity, and incast is spread in time, space, and priority rather than eliminated.
- **Telemetry saves GPU-hours**: PFC counters, link error rates, bandwidth baselines, and application throughput must be correlated across layers. Without that observability, a degraded transceiver or PFC storm silently converts a high-end fleet into a queue of waiting accelerators.

At fleet scale, the interconnect is part of the computer. Its value lies in how much useful computation it allows the participating accelerators to complete together. A fabric decision therefore has to be judged against the workload's communication pattern and synchronization schedule, because additional capacity helps only where those dependencies expose a communication limit.

What's Next: From wires to data pipelines

A fabric that completes every collective on time still leaves two flows of bytes unaccounted for. Each step needs its next input batch in place before the accelerators ask for it, and the fleet's training state must outlive the next hardware failure, which at this scale means writing multi-terabyte checkpoints on a schedule. Those checkpoint bursts are the traffic class this chapter isolated from the synchronization path, and for text pretraining they can outweigh the training corpus itself. Chapter 4 treats storage as the fleet's fuel line, sizing the tiers and pipelines that deliver inputs ahead of demand and persist model state without stalling the step.

4

Data Storage

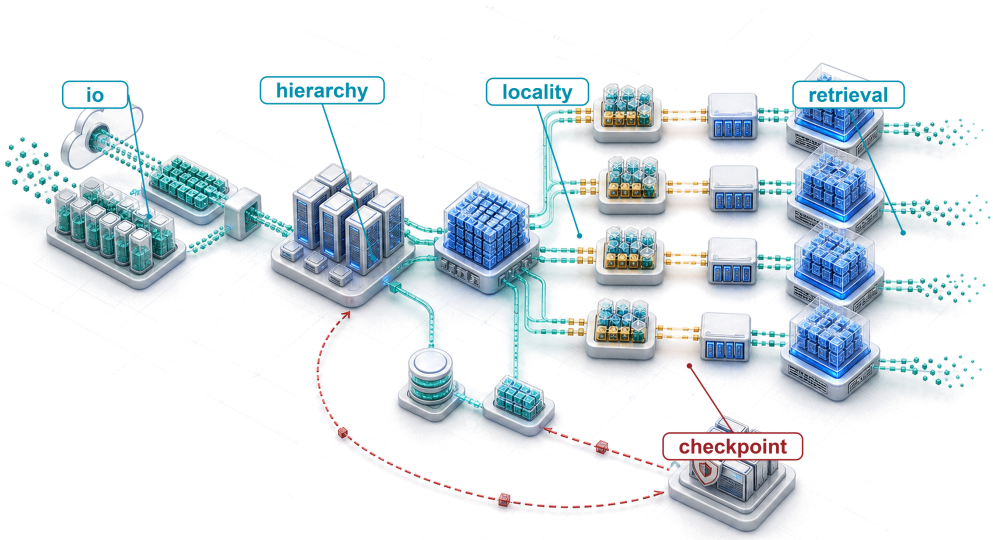

4.1 The Fuel Line
4.2 How ML Workloads Invert Storage Assumptions
4.3 The ML Storage Hierarchy
4.4 The Data Pipeline Equation
4.5 GPUDirect Storage and the CPU Data-Copy Bypass
4.6 Checkpoint Storage
4.7 Versioned Data and Commit Boundaries
4.8 Storage Economics
4.9 Storage for Serving
4.10 Storage Technologies That Move Tiers
4.11 Fallacies and Pitfalls
4.12 Summary

Purpose

Why does storage become the invisible bottleneck that prevents accelerators from reaching their potential?

A training job can leave expensive accelerators idle even when its models fit in memory and its network has ample capacity, because the next batch must arrive before computation needs it and a saved checkpoint must remain usable when a worker fails. Storage becomes an invisible bottleneck whenever accelerator benchmarks begin with data already in place, excluding the work needed to deliver inputs and preserve results. A pipeline that keeps one machine busy can stall once many workers compete for the same data source. A burst of checkpoint writes from one job can delay input reads for others, which turns a local pause into lost progress across the fleet. Adding compute increases demand without expanding the shared system's ability to serve it. How severely that contention hurts depends on the workload: compact training inputs may be easy to deliver even when saving model state dominates storage traffic, whereas larger samples can make continuous input delivery the limiting factor. Bringing data closer to the accelerator reduces those delivery delays, though it adds staging traffic, capacity cost, and the obligation to preserve complete versions, since a fast local save provides little protection if the only copy disappears with the failed machine. Storage decisions therefore govern both how quickly the fleet advances and how much of that progress survives disruption. In C^3 terms, storage design coordinates data movement and durable state so that available compute produces recoverable progress.

Learning Objectives

- Compare storage demands of full-corpus training, checkpoint writes, and model serving across local and shared tiers
- Calculate input bandwidth and mean cache service time from workload demand and tier behavior
- Design shard packaging, prefetching, locality, and direct-storage paths that keep fleet accelerators fed when small files or slow tiers would starve them
- Select checkpoint staging and replication strategies that balance pause time, recovery risk, and storage cost
- Specify storage commit and replay contracts that preserve correctness when distributed pipeline stages fail, including the provenance that generated training data must retain
- Analyze serving-time storage workloads, from bulk model loads and rollout bursts to retrieval indexes bound by random-access IOPS
- Evaluate staging and tiering choices by total delivery cost, including transfer fees and idle accelerators

4.1 The Fuel Line

The dense accelerator nodes of chapter 2 pack eight GPUs delivering petaFLOP/s of aggregate compute, and the InfiniBand fabrics of chapter 3 connect thousands of such nodes at hundreds of gigabits per second. Within the multi-tier fleet stack of figure 1.4, data storage completes that physical foundation, staging training data, model weights, optimizer state, and intermediate checkpoints at the right distance from the accelerator. An engine without fuel is expensive sculpture. The engineering question is how to deliver that fuel fast enough that the accelerators of a training cluster never starve.

The running cluster for that question is 256 eight-GPU nodes, 2,048 accelerators in all, deliberately twice the 1,024-GPU (128-node) reference cluster used elsewhere in this book, because the larger fleet makes the pressure on shared storage easier to see. Every worked number for the 175B job uses this running cluster; a side scenario with a different workload or fleet size names its own scale where it appears.

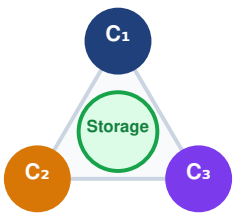

Storage is the physical substrate spanning Compute (C_1), Communication (C_2), and Coordination (C_3).

Consider the running example for the storage analysis, an instance of Archetype A (GPT-4/Llama-3). A 175-billion parameter language model trains on 1.5 trillion tokens of text: roughly 3 TB in compressed source form, or 6 TB once represented as 4-byte token IDs. That token budget is set upstream of storage. Under the compute-optimal scaling results of section 1.3, the training-token count grows in proportion to parameter count, so each step up the model-size curve enlarges both the corpus the fleet must hold and the bytes each epoch must read. This corpus already sits below the compute-optimal allocation for a 175B-parameter model, so a compute-optimal run would raise both figures further. The run makes a single pass (one epoch) over this corpus, reading every selected token once in an order determined by the shuffle seed. No small "hot" subset dominates access. Meanwhile, each training step consumes one global batch of 1,536 sequences (3,145,728 tokens), then waits for the next. On the running cluster the modeled step takes about 4.64 seconds, an estimate from the book's performance model rather than a measurement, so the single epoch lasts about 25.6 days. If storage cannot deliver that batch before the current step ends, the accelerators sit idle, and the organization pays for silicon that produces heat instead of gradients.

The problem is deceptive because storage technology has improved substantially. Non-Volatile Memory Express (NVMe) drives achieve 7 GB/s of sequential throughput, a figure that would have seemed out of reach a decade ago. An H100 moves data internally through high-bandwidth memory (HBM) at 3.35 TB/s, roughly 479× the sequential bandwidth of a single NVMe drive. This ratio contrasts two tiers; it is not the rate at which storage must supply training inputs, because HBM bandwidth also carries repeated weight, activation, and tensor traffic during computation. The storage feed requirement is instead determined by the input-batch bytes each accelerator needs per iteration. Dividing those bytes by the compute window gives one worker's required delivery rate; summing across workers gives the fleet demand. That rate must be evaluated at each tier using the

representation stored there, since compressed inputs and decoded tensors occupy different numbers of bytes. A hierarchy becomes useful when the durable source cannot meet that rate or its latency varies enough to interrupt computation. Faster tiers stage upcoming batches while the durable tier supplies capacity and preserves the source data.

The storage problem joins physics with economics. Physics favors the memory closest to the accelerator, which offers the highest bandwidth in the smallest capacity; economics favors shared storage, which pools far more capacity at lower cost per byte behind additional interfaces and service queues. Device technology, interconnects, and contention then determine the rate a particular consumer actually receives. The engineering task is to place and move data so that the accelerator receives each batch before it is needed. The resulting design problem is quantitative. It asks how fast each tier must be, how deep the pipeline must run, and which bytes are worth keeping close to the accelerator. Improving arithmetic throughput alone does not improve input delivery or checkpoint persistence. Those paths require their own bandwidth and latency budgets.

The canonical training-data footprint for this example is roughly 9 TB across the hierarchy, combining the example's compressed corpus and tokenized shards. Additional shuffled or packed variants can raise the staging footprint, but the per-epoch training read over 4-byte token IDs is 6 TB. This chapter adopts one resumable mixed-precision Adaptive Moment Estimation (Adam) checkpoint contract: FP16 model weights at 2 bytes per parameter, FP32 master weights at 4, and two FP32 Adam moments at 4 each, for 14 bytes per parameter. This is the 16-byte training state of section 2.5.3 without its 2-byte FP16 gradients, which are recomputed after restore, while scheduler, random-number-generator, and data-loader metadata are negligible at this scale. The resulting checkpoint is 2,450 GB every 10 minutes. The storage plan budgets a 30-day training window, which the modeled 25.6-day epoch fits inside. Over that window on the 256-node cluster, the storage system must deliver 6 TB of tokenized training data in the run's single epoch, absorb 10.6 PB of checkpoint writes, and stage model weights for evaluation runs. The sections that follow use this same contract for every checkpoint size, duration, retention total, and lifetime-write calculation.

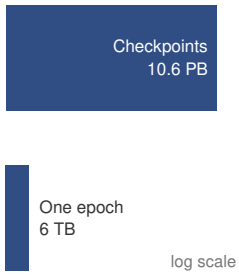

Over the run, checkpoint writes exceed one epoch of training-data reads by more than 1,000 times.

Those numbers force the storage path. ML access patterns first invert the assumptions behind conventional storage systems. That inversion forces a hierarchy from HBM to cold archive. The hierarchy then requires a pipeline equation that sizes delivery and direct data paths that remove host copies. It also requires a checkpoint protocol for the dominant write traffic, commit boundaries that keep every versioned artifact reproducible (generated training data included), and economics that decide which bytes belong at each tier. Serving-time model loading and retrieval indexes are variations on the same fuel-line problem, and devices that move the boundaries between tiers close the chapter. In each workload, the storage system must place the right representation at the right distance before the accelerator asks for it.

4.2 How ML Workloads Invert Storage Assumptions

A database administrator moving to an ML infrastructure team would find many familiar storage design principles inverted by the multi-terabyte sequential scans and checkpoint bursts of the 175B running example. Transactional databases emphasize small random accesses and strong consistency, issuing thousands of 4 KB reads per second to serve user queries. Storage designs for that regime combine caching algorithms that assume recently accessed records will soon be requested again, write-ahead logs that save each change before applying it, which supports point-in-time recovery, and redundant array of independent disks (RAID),[1] which trades capacity and performance for fault tolerance. Full-corpus pretraining breaks the locality and access-size assumptions behind that stack while retaining the need to protect durable state; fine-tuning, curriculum learning, hard-example mining, recommendation features, and retrieval introduce their own skewed working sets, so each inversion below is a property of the workload rather than of ML as a whole.

Full-corpus pretraining inverts these assumptions when samples are packaged into sequential shards (large files that each hold many samples), allowing the pipeline to stream through datasets that span hundreds of terabytes. Individual accesses are large (megabytes rather than kilobytes) because models consume batches of images or text sequences, while immutable, versioned training shards avoid read-write conflicts during a run even though checkpoint publication and live feature updates still require explicit consistency guarantees. Image and video workloads can demand hundreds of gigabytes per second for days or weeks, whereas compact tokenized text can require

[1] **RAID (redundant array of independent disks)**: A Berkeley taxonomy of drive-combining strategies that trade redundancy, capacity, and I/O performance (Patterson et al. 1988). RAID 0 stripes data without redundancy and can suit a disposable cache backed by a durable source. A drive failure then requires restaging; redundancy may still be worthwhile when that recovery cost or the required availability justifies it. Unreplicated checkpoints and sole dataset copies require a different durability contract.

much less input bandwidth even when checkpoint writes dominate total traffic. The mismatch between what storage was optimized for and what ML actually needs creates the **I/O wall** (principle 7): when storage throughput cannot deliver training data as fast as accelerators consume it, GPUs idle regardless of their computational power. For bandwidth-heavy image and video inputs, that wall arrives when accelerator count grows faster than provisioned storage throughput, making the *data pipeline*, not the *model*, the limiting factor whenever a fixed input service time cannot keep pace with faster compute.

↪ How storage stalls surface as communication or coordination bottlenecks in the C^3 taxonomy is developed in section B.4.6.

> **Systems Perspective 4.1: The moving I/O bottleneck**
>
> Hold the batch representation and storage path fixed while improving compute throughput. The batch now finishes sooner, leaving less time to deliver the next one. Prefetching can overlap delivery with execution, but it cannot overcome a sustained service-rate deficit. This compute-to-storage imbalance does not imply that every larger model is input bound. A model that performs more computation on the same batch can instead give storage more time. The design must follow workload service times, not advertised peak-throughput trends.

A normalized example isolates the effect. Suppose input delivery takes one quarter of the original compute time and overlaps perfectly with computation, so each batch takes the longer of the two service times. Doubling compute speed halves batch time, and quadrupling it halves batch time again, at which point compute and delivery take equal time. Beyond that point the pipeline is delivery bound. A sixteenfold faster accelerator still raises batch throughput only fourfold, and the remaining compute sits idle behind the input path. That plateau is an I/O wall for this workload and allocation. Faster storage, a more compact representation, or additional independent input paths can move it; deeper prefetching alone cannot.

The first inversion is access pattern. Where database workloads issue *random access patterns* that benefit from seek-time optimization, full-corpus pretraining performs large *sequential scans*. An epoch in the running example reads each selected sample once, in the order produced by shard and buffer shuffling, so the pattern resembles video streaming more than database queries. Storage systems optimized for random IOPS[2] waste capabilities on this workload, whereas systems optimized for sequential throughput excel. The penalty for the wrong pattern is large. A Gen4/Gen5 NVMe drive delivers roughly 7–14 GB/s for sequential reads, while small 4 KB random reads can fall near 0.5 GB/s, an order-of-magnitude to 30× drop. Hard drives pay an even steeper 100× throughput reduction because mechanical seek times dominate random access.

The shuffling that ML training requires adds a complication rooted in the requirements of stochastic gradient descent. Stochastic gradient descent requires that each mini-batch be drawn approximately uniformly from the training distribution; presenting samples in a correlated order (all samples from the same document, all images from the same class) biases the gradient estimates and slows or destabilizes convergence. True global shuffling would satisfy this requirement perfectly, but it requires random access across the entire dataset, destroying the sequential access pattern that storage hardware demands. Even for a multi-terabyte corpus, random per-sample seeks are I/O-prohibitive. A round illustrative dataset of 1 trillion 4-byte tokens (about 4 TB) stored as individual elements on NVMe and accessed at random would take hundreds to thousands of hours to read at the drive's random-read IOPS rate, rather than about ten minutes of sequential streaming. The practical solution is therefore a compromise that shuffles the order of large data shards and then shuffles samples within each shard's local buffer. This preserves sequential I/O, but adequate mixing depends on shard composition and buffer size, a dependence section 4.4.2 turns into a validation step for the loader.

The second inversion is working set size. Traditional applications rely on temporal locality. A web server repeatedly accesses popular pages, and a cache holding the top 10 percent of content serves 90 percent of requests. In the chapter's full-corpus pretraining scenario, the 6 TB serialized corpus has no "popular" token range; each selected token is consumed once per epoch. A cache of any practical size holds only a small fraction of the corpus, and most samples are effectively "cold" when accessed. This access pattern makes traditional recency caching much less effective. A least recently used (LRU)[3] cache can approach a 0 percent hit rate for a cyclic scan larger than the cache, because

[2] **IOPS (input/output operations per second)**: The metric that dominated storage procurement for decades because database workloads issue millions of small random reads. ML training inverts this priority: a pipeline streaming 256 MB shards sequentially needs sustained GB/s throughput, not per-operation speed. Provisioning storage by IOPS rating for an ML workload over-spends on random-access capability the pipeline never exercises.

[3] **LRU (least recently used)**: An eviction policy that discards the entry unused for the longest time, effective when recently accessed data is likely to be accessed again soon. Under independent reshuffling, some samples retained from the end of one epoch may be reused early in the next, so teams should measure that reuse before allocating DRAM that could otherwise serve as prefetch buffers.

each sample is evicted before its next access. With independently reshuffled epochs, cross-epoch hits are low but nonzero and depend on the cache fraction and sample order.

Other training policies create different locality. Multi-task or curriculum learning, hard-example mining, and filtered sampling access some subsets more frequently during particular phases. This temporary working set lets the NVMe tier exploit locality. Large-scale pretraining jobs that repeatedly scan a fixed corpus remain the scenario for the hierarchy derived here; systems serving another sampling policy should measure its reuse distribution before adopting the same cache design.

The third inversion is write pattern. Transactional systems generate continuous streams of small writes, each immediately durable. ML systems generate occasional large writes when saving checkpoints. A 175B parameter model checkpoint, including optimizer state, occupies roughly 2,450 GB. Saving it every 10 minutes generates concentrated bursts that saturate bandwidth for seconds, followed by long idle periods. These synchronized bursts, called checkpoint storms and defined formally in definition 4.5, must be absorbed by the shared file system without disrupting ongoing training reads. The bursty write pattern is particularly challenging because all nodes in the cluster write their checkpoint shards simultaneously. If 256 nodes each write 10 GB, the parallel file system receives 2.4 TB of writes in a single burst, which blocks training if the job writes synchronously to that tier. Local staging and asynchronous replication can shorten this pause, as section 4.6 develops. When checkpoint and gradient traffic share one fabric, traffic isolation also keeps the burst from delaying the job's collectives.

The fourth inversion is subtler. The read/write ratio shifts by lifecycle phase. During the intervals between checkpoints, the input pipeline reads continuously. Across a complete run, however, the read/write ratio depends on input size, epoch count, checkpoint size, and save frequency. The running language-model example produces far more checkpoint bytes than training-input bytes despite spending most of its time between saves. During data preprocessing, both reads and writes are heavy, and the access pattern more closely resembles a MapReduce job than a training loop. Tokenization, deduplication, and shuffling scan the raw corpus once (large sequential reads) and write intermediate artifacts (large sequential writes). The next stage then scans those artifacts again. A 3 TB compressed text corpus typically expands to 6 TB of tokenized shards before the first training epoch begins, loading each storage tier differently from the read-only training phase that follows. No single storage configuration optimizes for all phases, which is why ML systems require a multi-tier hierarchy rather than a single storage technology.

↩ Traffic isolation and quality-of-service mechanisms are established in section 3.8.2.

A fifth inversion emerges when comparing training and inference workloads. *Training* reads datasets sequentially and writes checkpoints in bursts. Weight-resident *inference*, by contrast, performs one large model read at startup and then serves requests from HBM. Its storage challenge is cold-start latency: the time required to load a model when scaling up or recovering from failure. The limiting constraint flips from sustained dataset throughput to a bulk read whose duration depends on which tier supplies the weights. Retrieval, KV-cache offload (moving a request's saved attention state out of HBM), and models larger than aggregate HBM create additional serving-time I/O, so this weight-resident case is a baseline rather than a universal inference pattern. Section 4.9 derives the load times and connects that baseline to the retrieval workload that follows.

Serving at fleet scale also creates a fan-out read burst. A serving cluster with 100 replicas of the 175B running model holds 35 TB of FP16 weights, and a model rollout must deliver a new version to every replica within minutes while the old version keeps serving. The burst is the checkpoint storm in reverse. Instead of many nodes *writing* to shared storage at once, many nodes *read* the same data at once.

Table 4.1: ML Workloads Invert Traditional Storage Assumptions: The five inversions of full-corpus training and serving. Where databases optimize for random IOPS with cacheable working sets, ML training streams sequentially through datasets that exceed all cache levels and writes state in bursts, its read/write balance shifts by lifecycle phase, and serving replaces dataset streaming with bulk model loads.

Workload pattern	Traditional assumption	ML reality
Access pattern	Random access	Sequential streaming
Working set	Fits in cache, with strong temporal locality	Exceeds all cache levels; low locality in full-corpus scans

Workload pattern	Traditional assumption	ML reality
Write pattern	Continuous small writes	Bursty large writes
Read/write ratio	Balanced	Phase-dependent (read-only to write-dominated)
Training vs. inference	One access pattern for all phases	Corpus streaming in training; bulk model loads and fan-out in serving

Together, the five inversions collected in table 4.1 reduce to one procurement rule, which is to optimize for streaming reads and bursty writes rather than database-style random IOPS. An organization provisioning storage for ML based on database-era heuristics will over-invest in random IOPS for sequential training shards, under-invest in the sustained bandwidth those shards require, and fail to account for the bursty write patterns that checkpointing creates.

Checkpoint 4.1: Storage workload analysis

Consider the storage subsystem for a new ML training cluster with 512 GPUs. The primary workload will train large language models on a 10 TB text dataset.

- ☐ **Database heuristics**: use the five storage workload inversions to name at least three database optimization strategies that would be counterproductive for this workload.
- ☐ **Epoch reads**: calculate the total data volume read when each training epoch reads the full 10 TB dataset once and the cluster trains for 100 epochs, then compare it to a typical database workload.
- ☐ **Checkpoint bandwidth**: calculate the minimum write bandwidth required when the model saves a 500 GB checkpoint every 15 minutes and each save must complete within 30 seconds.
- ☐ **Write duty cycle**: calculate the percentage of time the storage system spends in "write mode" vs. "read mode" under bursty checkpoints every 15 minutes.

Translating these five inversions into a concrete architecture requires separating how a dataset is packaged from how fast each storage tier must serve it. The required delivery rate depends on workload demand rather than parameter count alone. It equals the aggregate input bytes consumed during a delivery window divided by that window, a rate that representation, batch size, parallelism, and compute time jointly set (section 4.4) and that prefetching alone cannot sustain when demand exceeds a tier's allocated share. In the running example, the 1.5 trillion tokens of training data produce roughly 6 TB of sequential token-ID reads per epoch from a 3 TB compressed source corpus. Stored as one file per sample, that corpus would spend its storage budget on file opens rather than payload bytes, the metadata bottleneck that section 4.3.6 develops as the small file problem. Preprocessing the corpus into large sequential shards (Aizman et al. 2019) amortizes open, seek, and close costs across many tokens without changing the payload byte rate (section 4.3.5). Where those shards should live, and how close to the accelerator, is the question the storage hierarchy answers.

4.3 The ML Storage Hierarchy

No single storage device can hold a petabyte-scale training corpus, absorb multi-terabyte checkpoints, and still feed an accelerator at terabytes per second. HBM supplies high local bandwidth at limited capacity, whereas object storage pools durable capacity behind a much slower service and client path. A multi-tier hierarchy combines these resources rather than requiring one technology to meet every demand. Each tier offers a different balance of bandwidth, latency, capacity, and cost. The hierarchy extends the familiar processor memory hierarchy across nodes. The new fleet-level decision is how local staging reduces demand on shared storage while preserving a durable, consistently versioned source. Comparing the tiers fairly requires stating each resource at its own scope, per accelerator, per node, or as an aggregate shared service (table 4.2).

Table 4.2: Extended Memory Hierarchy for ML Systems: HBM bandwidth is per accelerator, host DRAM and local NVMe are per node, and parallel-file-system and object-storage figures are aggregate service rates. These different scopes are not directly comparable tier ratios; divide shared bandwidth among active clients and account for each client path before sizing prefetching and caching. Memory costs are acquisition dollars per GB, while storage costs are recurring dollars per GB-month; comparing them requires an explicit amortization period and utilization assumption.

Storage tier	Typical capacity	Bandwidth	Latency	Illustrative cost basis
GPU HBM	80 GB	3.35 TB/s	~273 ns	~$15/GB acquisition cost
Host DRAM	512 GB–2 TB	200 GB/s	~100 ns	~$3/GB acquisition cost
Local NVMe SSD	4–30 TB	7–25 GB/s	~100 µs	~$0.10/GB/month provisioned
Parallel File System	100+ PB	1+ TB/s aggregate	~1 ms	~$0.03/GB/month provisioned/service
Object Storage	Very large pool	100 GB/s aggregate	~50 ms	~$0.02/GB/month service
Archive/Cold Storage	Very large pool	1 GB/s	Minutes to hours	~$0.004/GB/month service

HBM's advantage over host DRAM in table 4.2 is *bandwidth*, not *latency*, because its lead comes from wide stacked interfaces rather than faster cells (section 2.3.1). The latency cliff between adjacent tiers opens only at NVMe and below, where each access adds media time and storage-protocol processing and, beyond the node, network hops and queueing. The storage hierarchy model (principle 8) therefore governs storage tiering and pipeline design. The local tiers exchange bandwidth for capacity, while shared tiers aggregate many devices and expose a different client path. Because per-client bandwidth depends on allocation and contention rather than a fixed factor between tiers, data format choices, caching strategies, prefetch buffer sizing, and tiering policies all exist to manage the movement of data upward through the hierarchy so that the accelerator never starves.

The rest of the chapter numbers the six tiers by distance from the accelerator, from Tier 0 (HBM) through Tier 1 (host DRAM), Tier 2 (local NVMe), Tier 3 (parallel file system), and Tier 4 (object storage) to Tier 5 (archive). Distance alone hides an important distinction, because each tier's bandwidth figure describes a different scope: per accelerator, per node, or a shared service divided among its clients (figure 4.1).

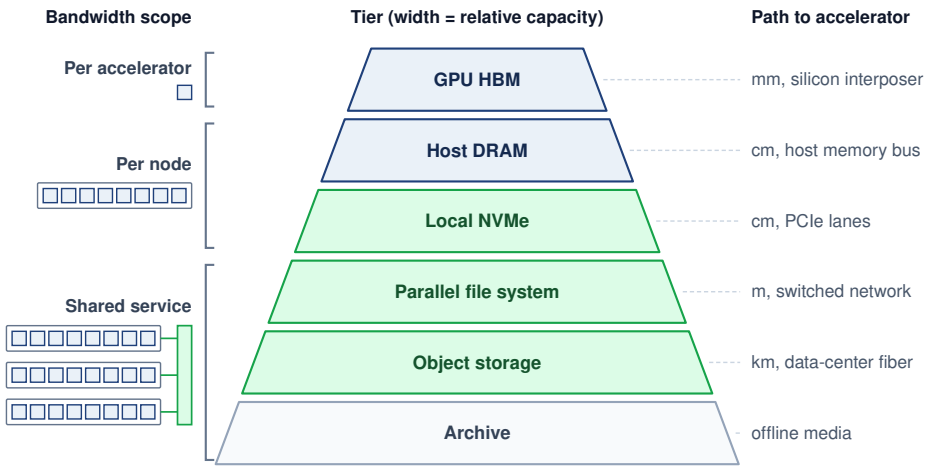

Figure 4.1: ML Storage Hierarchy: The six tiers of table 4.2 stacked by distance from the accelerator, with band width showing schematically how capacity grows toward the base; the table carries the numbers. The left brackets give the scope of each tier's bandwidth figure, and each glyph shows what shares it: one accelerator, the accelerators of one node, or many nodes drawing on one aggregate shared service. The right column gives the physical path each byte travels to reach the accelerator. Blue memory tiers are bought per GB; green storage tiers are rented or provisioned per GB-month; the gray archive tier is offline media.

The physical tiers of figure 4.1 govern the trade-off among access distance, capacity, and cost. Shared services can aggregate more bandwidth than an individual local device, so the ordering is not a monotonic ranking of aggregate throughput. This trade-off is not arbitrary; it reflects the physics of data proximity. HBM sits on the same silicon interposer as the accelerator, connected by thousands of parallel traces measured in millimeters, whereas host DRAM connects to the CPU memory controller and reaches the accelerator across the host-device interconnect, and local NVMe reaches across the circuit board over PCIe lanes. Beyond the node, the parallel file system traverses meters of cable and network switches, while object storage may span kilometers of fiber between data centers. Distance, protocol overhead, and contention can increase access latency, while media technology and provisioning determine capacity cost. Distance is therefore a guide rather than a measurement; only the complete client path establishes all three properties for a given consumer.

The engineering challenge is to deliver each batch to HBM before the accelerator needs it. In the running example, the durable corpus lives in object storage (Tier 4) as a 3 TB compressed source copy plus 6 TB of tokenized training shards, but the accelerator needs each batch in HBM (Tier 0) before the current step ends. The chapter's reference design fixes one path between them. At job start, each node copies its share of the tokenized shards from object storage to its local NVMe (Tier 2). The single training epoch then reads from that local copy, prepares batches in host DRAM, and transfers them to HBM. Local reads keep slow object-store responses out of each training step, the reason section 4.8 gives for staging a corpus this job reads only once. The parallel file system (Tier 3) holds checkpoints rather than training inputs. Streaming the shards directly from object storage, or serving them from a parallel file system, are the alternatives the chapter weighs against this design, and the hierarchy describes placement choices, not a requirement that every byte visit every level.

4.3.1 Bandwidth cliffs between tiers

Stepping down one level from a single accelerator's HBM (3.35 TB/s) to the whole host's DRAM (~200 GB/s) drops bandwidth by roughly 16.8×, and that comparison still understates the cliff on an eight-GPU node because all eight accelerators share the same host-memory bandwidth. Comparing adjacent tiers requires making the scope of each rate explicit (per accelerator, per node, or across a shared service), as bracketed in figure 4.1. Between host DRAM and NVMe, the ratio is roughly 7.1× (200 GB/s vs. ~28 GB/s from a 4-drive RAID-0). Between NVMe and a parallel file system, the ratio depends on the per-node allocation. If a 1 TB/s aggregate parallel file system (PFS) serves 256 nodes, each node receives roughly 4 GB/s, a 7× reduction from the local RAID-0 stripe. Object-storage throughput likewise depends on the client path and allocation; there is no fixed ratio to a parallel file system.

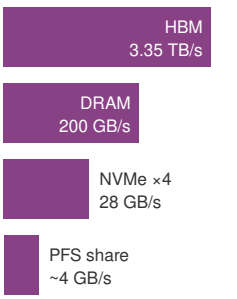

Per-client bandwidth falls roughly 7 to 17× at each step from HBM to a node's share of the parallel file system.

As section 4.1 established, a cliff is not a feed requirement; the feed rate remains input-batch bytes divided by iteration time, summed over distinct consumers. What the cliffs decide is where buffering must sit. Along the chosen path, host DRAM can buffer prepared batches and NVMe can cache shards from shared storage. Each buffer must cover the delay of the tier below it while that tier sustains the consumption rate; durable repositories need not act as intermediate buffers.

The same hierarchy can be analyzed as a cache rather than as six isolated devices. If a request hits the hot tier with probability h, taking T_{hot}, and otherwise falls through to a colder tier, taking total end-to-end miss-path time T_{cold} (including the failed hot-tier lookup), its mean service time is

$$\mathbb{E}[T] = hT_{\text{hot}} + (1-h)T_{\text{cold}}. \tag{4.1}$$

For fixed-size records served serially, the corresponding effective bandwidth is the record size divided by $\mathbb{E}[T]$. The arithmetic makes small miss rates consequential. With a 95 percent hit rate and a cold-tier access 100× slower than a hot-tier access, $\mathbb{E}[T] = 5.95T_{\text{hot}}$, so the effective request rate is only about one-sixth of the all-hot rate. Concurrent requests and prefetching can overlap service, so this mean is a sizing baseline rather than a fleet-throughput or latency guarantee. Correlated misses or a heavy-tailed cold tier can drain the prefetch buffer even when the average satisfies the pipeline budget, so production designs must check a tail percentile and miss bursts as well as $\mathbb{E}[T]$.

The batch-volume arithmetic also explains why increasing cluster size creates storage pressure. Aggregate demand is each consumer's batch bytes divided by its compute window, summed over consumers, so at a fixed per-node rate it grows linearly with node count. The per-node rate itself

spans orders of magnitude across data modalities, as text and image training on the running cluster show.

Napkin Math 4.1: Text vs. image bandwidth

Problem: How much aggregate storage bandwidth does a 2,048-GPU cluster require for text training vs. image training?

Setup: Compare two declared batch representations at the same compute window and number of independent input consumers. The result depends on batch volume and rate, not modality alone. Both jobs are generic ones sized to the running cluster, not the 175B running job. Each GPU consumes one per-GPU batch every step, a pace that fits a model far smaller than 175B parameters.

For text training, the demand is low. With a per-GPU batch of one packed sequence of 4,096 tokens and a 200 ms step time, the aggregate bandwidth is:

$$2{,}048 \text{ GPUs} \times 4{,}096 \text{ tokens/GPU} \times 4 \text{ bytes/token} \div 0.2\text{s} \approx 167.8 MB/s$$

This is easily served by a single network-attached storage node, one file server that clients reach over the network.

Image training demands far more. Using a common batch size of 256 per GPU (ImageNet at 224×224, roughly 116 KB/image), the aggregate bandwidth is:

$$2{,}048 \text{ GPUs} \times 256 \text{ images/GPU} \times 115.7 \text{ KB/image} \div 0.2\text{s} \approx 303.2 GB/s$$

Systems insight: Under the stated batch representations and compute times, image training requires 1,800× higher input bandwidth than text training. The text example is *volume-heavy* but *bandwidth-light*. One network-attached storage server can carry its input feed, so its storage design is shaped by total corpus volume and, as later sections show, by checkpoint traffic. The image example is *bandwidth-heavy*, with a feed that exceeds any single storage server and calls for the aggregate bandwidth of many servers or of local caches. Whether either demand becomes an actual bottleneck still depends on the measured delivery rate of the path that serves it.

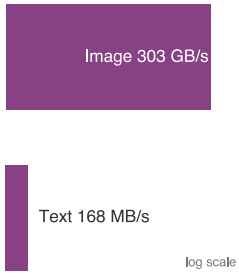

Image training demands 1,800× the storage bandwidth of text.

The 175B running job sits lower still. Its 6 TB epoch spread over the modeled 25.6-day run averages about 2.7 MB/s for the whole cluster, corpus bytes divided by run length, roughly 60× below the notebook's generic text job. Each token of a 175B model carries so much computation that the corpus trickles in, which is why the running example's storage design is shaped by checkpoint traffic rather than by its input feed.

The image result is a floor for visual data rather than a ceiling. At about 303.2 GB/s across 256 nodes, the image job already consumes nearly a third of a 1 TB/s parallel file system built from 100 data-serving nodes (the architecture section 4.3.6 develops). Higher-resolution images and video clips, at megabytes per sample rather than the roughly 116 KB of a compressed ImageNet image, raise the per-node demand by an order of magnitude or more, and a fleet several times larger multiplies it again. Demand at that scale exceeds what a shared file system of this size can serve, so the architecture must rely on local NVMe caching to take repeated reads off shared storage.

The bandwidth cliff between tiers also has implications for the data format at each level. Computation consumes float16 tensors, packed token IDs, or other supported representations, while NVMe and object storage can retain compressed records. Compression reduces bytes stored and transferred but adds decoding work. The chosen format must let CPU or GPU decoding keep pace with consumption; maximum compression can reduce end-to-end throughput when decoding becomes the limiting stage. A host-decoded path stages the resulting tensors in DRAM, while a GPU-decoded path moves the transformation and its resource demand to the accelerator.

The format challenge intensifies for multi-modal training, which combines text, images, audio, and video in a single model. Serialized sizes span orders of magnitude, from 4 bytes per text token to roughly 116 KB per compressed image to megabytes per short video clip, depending on

representation, compression, resolution, and duration, and each modality requires its own decoding and augmentation work. A multi-modal training job must manage these parallel streams and provision for their combined service demand. For a corpus containing 3 TB of text, 50 TB of images, and 200 TB of video, video dominates retained capacity. Its bandwidth contribution also depends on how frequently training samples it and how many bytes each sampled batch reads. Corpus size alone determines neither the delivery rate nor the modality's contribution to model quality.

Tracing how each tier meets these bandwidth and format constraints begins at the top of the pyramid and steps outward from the accelerator to the archive. That descent starts where computation actually consumes bytes: HBM, the tier fast enough to keep accelerator arithmetic fed but far too scarce to hold the full training corpus.

4.3.2 Tier 0 GPU HBM

The scarcest resource in the system sits closest to the arithmetic. On the H100 used here, HBM holds the active weights and activations served to the accelerator, while caches and on-chip buffers reuse portions of that state. Its stacked packaging supplies 3.35 TB/s of aggregate bandwidth, the top of the cliffs just measured. That bandwidth is what makes large-scale deep learning feasible, because a matrix multiplication involving billions of parameters reads those parameters from memory on every forward and backward pass. What HBM lacks is room.

↩ High-bandwidth memory packaging, the 3D stacking and through-silicon vias behind this bandwidth, is established in section 2.3.1.

For placing the unpartitioned running model, the immediate constraint is *capacity*. Once state fits, HBM bandwidth can still limit execution. The node-memory budget of section 2.5.3 already shows how far one 80 GB device falls short of the 175B running example. The FP16 weights alone, 350 GB, need at least five H100 accelerators. The mixed-precision contract this chapter checkpoints also retains 700 GB of FP32 master weights, so weights, gradients, master weights, and Adam moments occupy 2.80 TB before activations and communication buffers. State that does not fit in one HBM pool must be split across accelerators or kept outside HBM and fetched when needed. Splitting it across the cluster makes the run possible, but every shard becomes a separate piece of state that lower tiers must stage, save, and restore in step with the others. HBM scarcity is what forces the rest of the hierarchy to exist, and HBM is the destination that every lower tier exists to serve.

Because HBM capacity is so limited relative to both model size and dataset size, the accelerator processes data in batches. Each batch occupies a fraction of HBM for the duration of one forward-backward pass, then is discarded to make room for the next. The rate at which batches must be supplied sets the bandwidth requirement for all lower tiers.

The razor-thin margin is a defining feature of large model training, visible in the HBM budget that each 80 GB H100 carries while the fleet trains the 175B parameter model (table 4.3). The table treats partitioning as a storage layout, and the chapter keeps this one layout throughout. Tensor parallelism splits each layer eight ways across the GPUs of one node, and pipeline parallelism places successive groups of layers on eight nodes, so one complete replica of the model and its training state spans sixty-four GPUs, each holding a 1/64 slice. Data parallelism repeats that replica thirty-two times across the running cluster, and every replica holds an identical copy of the state. Sharding the replicated optimizer and gradient state across data-parallel replicas is the second stage of the Zero Redundancy Optimizer (ZeRO), a training technique in its own right; its third stage also partitions the weights. This layout needs neither, because the unsharded slices already fit. Here only the resulting per-GPU slices matter, because they set the storage pressure HBM sees.

↪ Sharded data parallelism (ZeRO and FSDP), which partitions optimizer, gradient, and weight state across workers, is developed in section 6.3.3.

Table 4.3: HBM Memory Budget for 175B-Parameter Training: Per-GPU allocation when one replica of a 175B model and its unsharded mixed-precision Adam state is split across sixty-four GPUs, eight ways inside each node and across a pipeline of eight nodes. The weight, optimizer, master-weight, and gradient rows are each GPU's 1/64 slice. The resulting budget occupies most of an 80 GB HBM device, leaving limited buffer for the incoming data pipeline. Any delay in fetching the next batch from host memory risks starving the accelerator.

Component	Size per GPU
Model Weights (FP16, replica slice)	5.5 GB
Optimizer State (replica slice)	21.9 GB
FP32 Master Weights (replica slice)	10.9 GB
Activations (variable by sequence length)	10–20 GB

Component	Size per GPU
Gradient Buffers (FP16, replica slice)	5.5 GB
Communication Buffers (NCCL)	2–4 GB
Total Occupied	**55.8–67.8 GB**

Even with aggressive partitioning, the total memory footprint reaches 55.8–67.8 GB in decimal gigabytes, about 64.9–78.9 percent of the HBM pool (the device's nominal 80 GB is a binary capacity, slightly larger than 80 decimal gigabytes). The remaining headroom shrinks further once fragmentation and framework workspaces are included, and that headroom is all the incoming data pipeline has to work with. An input batch and its activations occupy HBM for one training step, while model weights and optimizer state persist across steps. Prefetching the next batch therefore competes with the live training footprint for scarce capacity; a deeper input queue is useful only if the memory budget can accommodate it.

From the data pipeline's perspective, Tier 0 is not a storage tier to be managed but a constraint to be satisfied. The pipeline's purpose is to ensure that the batch the accelerator needs *next* is already resident in HBM before the *current* batch's computation completes. If it arrives late, the accelerator stalls. If it arrives early, it consumes HBM that could hold activations. The tension between "just in time" and "just too late" defines the pipeline's buffer management strategy, quantified in section 4.4.

4.3.3 Tier 1 host DRAM

One level below HBM, host DRAM serves as the staging area for the data pipeline. Every byte of training data that reaches the accelerator passes through host DRAM first (unless GPUDirect Storage bypasses it, as described in section 4.5). A typical training node contains 512 GB to 2 TB of system memory shared across the host CPU and its peripherals. The bandwidth between host DRAM and the accelerator is limited to what PCIe Gen 5 (about 64 GB/s per direction, 128 GB/s bidirectional) can provide. A faster path requires a CPU-accelerator superchip, a module that joins the two processors over a coherent NVLink-class host link, which approaches the 900 GB/s that NVLink provides between accelerators. Host DRAM plays three roles in the ML storage hierarchy. It is the prefetch buffer, a home for embedding tables too large for HBM, and the augmentation workspace. The data loader occupies two of these roles at once, because its augmentation stage and the prefetch buffer it fills both live in host DRAM.

The data loader pipeline that runs in host DRAM works as it does on a single node. Reader threads fetch compressed samples, workers decode and augment them, and a collation stage assembles batches in pinned, or page-locked, memory for the host-to-accelerator direct memory access (DMA) transfer. At fleet scale, these stages compete with the node's other demands for the same host DRAM.

Host DRAM's most critical function is serving as a prefetch buffer. When the accelerator finishes processing batch k, batch $k+1$ should already be assembled in host memory, ready for transfer to HBM. The depth of this prefetch buffer determines how much I/O variance the pipeline can absorb without stalling.

↪ The single-node overlap techniques behind this loader (prefetching, pinned host buffers, and double buffering) are collected for reference in section A.2.1.

Recommendation workloads, the regime of Archetype B (DLRM at Scale), place a different demand on host DRAM: hosting embedding tables that can exceed 100 GB, far too large for HBM. These tables reside in host DRAM and are accessed through lookups that fetch only the rows needed for the current batch. The bandwidth between host DRAM and HBM becomes the critical bottleneck for these workloads, and some systems go further, pooling embedding tables in the DRAM of dedicated CPU servers and returning only the requested rows over remote direct memory access (RDMA).

Host DRAM also provides the augmentation workspace that the CPU pipeline requires. Data augmentation operations (resizing images, applying noise) execute on the CPU and require temporary memory for intermediate results. A training pipeline that applies five augmentations to a 256-image batch at 116 KB per image needs tens of megabytes of working space for each augmentation stage. Although modest per-batch, this memory accumulates when multiple data loader workers run in parallel. GPU-side decoding with libraries such as the NVIDIA Data Loading Library (DALI) moves decode and augmentation onto the accelerator, so smaller compressed bytes cross the host link and host DRAM bandwidth demand falls. The work then consumes HBM capacity and compute cycles instead, which helps an input-bound job and hurts a compute-bound one.

↪ Sharding embedding tables across servers, and the sparse row exchanges that result, is developed in section 6.5.6.

The three roles interact in subtle ways. The prefetch buffer and the augmentation workspace compete for the same physical DRAM, and embedding tables consume capacity that could otherwise serve as deeper prefetch queues. A node with 512 GB of DRAM hosting a 200 GB embedding table has only 312 GB remaining for prefetching and augmentation. If the data loader uses 8 workers, each maintaining a decode buffer of 1 GB, the effective prefetch capacity drops further. System architects must balance these competing demands by profiling the actual memory consumption of each pipeline stage and provisioning DRAM accordingly.

The physical layout of DRAM also matters. On a dual-socket node, each socket's DRAM forms a non-uniform memory access (NUMA) domain local to the GPUs on that socket's PCIe root complex, and reaching the other socket's DRAM crosses the inter-socket link, with a penalty that depends on the processor, topology, and concurrent traffic (section 2.5.1.1). NUMA-aware allocation keeps each data loader worker and its prefetch buffer in the same NUMA domain as the GPUs it serves; the gain should be measured on the target node rather than assumed.

Host DRAM sits below the first cliff of section 4.3.1, and the host-to-GPU interconnect adds a separate cliff, with effective transfer bandwidth depending on whether the node uses PCIe or NVLink. In the conventional path that stages payloads through host memory, a depleted DRAM prefetch buffer leaves the accelerator waiting for lower tiers. The direct path of section 4.5 removes that staging copy when the data layout and devices support it. In the running example, the 256-node cluster requires each node's host DRAM to sustain a continuous flow of packed token batches ready for PCIe transfer. If the NVMe-to-DRAM read pipeline falls behind by more time than the prefetch buffer covers, the buffer drains and the accelerator idles until the next batch arrives.

4.3.4 Tier 2 local NVMe

When the working set exceeds host DRAM, the system falls to Tier 2: the local NVMe drives attached directly to the compute node. NVMe[4] provides a high-performance protocol designed specifically for solid-state drives, achieving 7 GB/s of sequential throughput per drive. NVMe's deep queues and PCIe lanes, rather than faster flash, explain why it delivers 10× to 50× the throughput of SATA SSDs (capped near 550 MB/s by the link) with identical NAND flash, provided the loader keeps many asynchronous reads outstanding. With four drives in a RAID-0 configuration, a single node can sustain roughly 28 GB/s of sequential reads before overhead, sufficient to stream the 6 TB serialized dataset from local disk in about 3.6 minutes under ideal sequential conditions.

[4] **NVMe (Non-Volatile Memory Express)**: A storage protocol designed for SSDs, connecting directly to the CPU via PCIe lanes. Its many deep command queues, replacing the single 32-command queue of the older SATA interface, are what allow a multi-worker ML data loader to keep enough requests outstanding to saturate the drive's bandwidth.

In ML training, local NVMe acts as a cache, storing data shards fetched from distributed storage. Workers can reread resident shards across epochs without fetching them over the network again. The savings depend on reuse, cache capacity, and whether remote reads cross a charged transfer boundary (section 4.8). Staging pays an initial transfer cost and provisions local capacity in exchange for avoiding later remote reads; it does not eliminate storage cost or read latency.

The local cache pattern requires careful capacity planning. A training node with four 7.68 TB NVMe drives provides approximately 30.7 TB of local storage. For the running example, the 6 TB serialized dataset fits comfortably on a single node's local storage with room for checkpoint staging and temporary augmentation buffers, and each node holds only its assigned share. A multi-modal training job in which any node may read any sample combines 20 TB of images, 10 TB of text, and 5 TB of audio for a total of 35 TB, exceeding one node's local capacity by about 4.3 TB and forcing the pipeline to stream from the parallel file system for at least part of the dataset. The design trade-off is between provisioning more local NVMe (which increases node cost) and accepting network-dependent reads (which risks latency spikes).

The I/O wall takes a deceptive form at this tier. The HBM-to-NVMe ratio of section 4.3.1 does not set demand, so the question here is whether the stripe meets the input-batch rate, not whether it approaches HBM. Pipelining overlaps that delivery with compute (section 4.4), while GPUDirect Storage (section 4.5) can remove host payload copies when those copies limit throughput. The wall is easy to miss because NVMe performance is excellent by historical standards. A storage engineer accustomed to HDD-era throughput of 150 MB/s might view 28 GB/s from a local RAID-0 stripe as more than enough. Whether 28 GB/s suffices depends on the batch representation and consumption rate. Compact text may fit within this budget, while image or video inputs may require more bandwidth; overlap hides transfer time only when sustained supply keeps pace with demand.

Local NVMe also supports local checkpoint staging. Writing a consistent snapshot locally and replicating it to shared storage in the background can shorten T_{write}, the time training pauses to save a checkpoint, when the shared write would otherwise be slower. Resuming after a local save accepts a period during which a node failure can destroy the new shard. The storage design must therefore specify both the exposed pause and the later durability boundary, while retaining a previous recoverable generation, one complete saved version of the checkpoint. Section 4.6 develops this trade-off and its timing model.

The same local NVMe tier also pays for itself by recovering utilization lost to remote-read latency. When congestion on a shared file system drains the host prefetch buffer and exposes I/O wait on every step, staging shards onto node-local flash converts that exposed remote wait into reclaimed accelerator compute time.

Napkin Math 4.2: Return on investment of local NVMe caching

Problem: A vision-model training pipeline, a separate workload from the 175B language model, runs each step in 800 ms on the 1,024-GPU reference cluster. It reads from a shared parallel file system, and even with prefetching, congestion bursts on that shared service drain the loader's buffer and leave an average of 150 ms of exposed I/O wait per step. Assume that serving the same shards from local NVMe leaves only 10 ms of exposed wait per step, the residue of local tail latency that the prefetch queue does not cover. How much does adding local NVMe SSDs to each node improve GPU utilization?

Math: GPU utilization here means the fraction of step time spent in computation, an active-time fraction rather than model FLOPs utilization.

1. **Remote Only**: 800 ms / (800 ms + 150 ms) $\approx$ 84.2 percent.
2. **Local Cache**: Prefetching from local NVMe removes the congestion bursts of the shared service, leaving the assumed local residue.
   - New Util: 800 ms / (800 ms + 10 ms) $\approx$ 98.8 percent.

Systems insight: Local storage can act as a *GPU utilization multiplier*. In this scenario, adding a $500 NVMe drive to a node whose accelerators each cost $30,000 recovers 14.6 percent of the GPU's capacity that was previously wasted on I/O wait. Across the 1,024-GPU cluster, that utilization gain is equivalent to adding 149 GPUs worth of useful work without buying that many more accelerators. The gain is exactly the exposed remote wait that local reads remove, so it is large only when shared-storage congestion leaves a wait of this size; where it does, local NVMe is not auxiliary but the physical buffer that decouples expensive compute from unpredictable shared storage.

A practical concern at this tier is SSD endurance. NAND flash memory can sustain a limited number of write-erase cycles before the cells degrade. Enterprise NVMe drives are rated for 1 to 3 drive writes per day (DWPD) over a 5-year lifespan. For a 7.68 TB drive at 1 DWPD, this means the drive can absorb 7.68 TB of writes per day, or roughly 14 PB total over its lifetime.

ML training workloads are predominantly read-heavy (the dataset is written once and read many times), which is favorable for SSD endurance. However, checkpoint writes can be intensive. If each node saves its roughly 10 GB share of the running example's checkpoint every 10 minutes, that is about 1.4 TB per day of checkpoint writes spread across the node's drives, well within a four-drive node's 1 DWPD budget. The risk emerges when local NVMe is used as a staging buffer for both checkpoint writes and dataset caching. The combined write volume from initial dataset staging plus repeated checkpoint saves must remain within the drive's endurance rating. Within that rating, read-heavy training fleets rarely wear out their drives, so a storage refresh cycle is usually driven by density rather than wear, and its timing belongs to the lifecycle cost analysis of section 12.9.3.

4.3.5 Data formats for training

Local NVMe can stream 28 GB/s per node, but packaging determines how much of that bandwidth reaches the loader. The same multi-terabyte dataset stored in different formats can yield I/O

throughput that differs by over 100× on identical hardware, because each format charges a different access cost: per-file metadata, unnecessary feature reads, or sequential shard traversal. A directory containing one JPEG per sample pays a separate file open and metadata lookup for every image. At roughly 10 to 50 microseconds per open in the kernel's file-system layer, a node reading 8,000 images per second spends 80 to 400 ms of CPU time per second on opens alone, before it decodes a single byte. On a shared file system the same operations also queue at a metadata service, the fleet-scale form of this cost that section 4.3.6 develops as the small file problem. A standard ImageNet training pipeline makes the cost of the wrong format visible.

Napkin Math 4.3: The ImageNet bottleneck analysis

Problem: A ResNet-50 training job on ImageNet (1.28M, ~116 KB average) targets 1000 images/s. Can individual JPEG files on an HDD sustain that rate, or does the job need NVMe?

Math:

1. **Raw Bandwidth**: 1000 images/s × 116 KB = 116 MB/s.
2. **HDD reality**: A 7200 RPM hard disk drive (HDD) delivers 100 IOPS at random reads. Sustaining 1000 images/s requires 10× more IOPS than the disk provides.
3. **Result**: When every image is a separate random read, shuffling individual files on an HDD will starve the GPU, capping utilization at about 10 percent before decode, seek, and file-system overheads.

Systems insight: Two fixes work, and they remove different limits. NVMe answers a random read in roughly 100 µs rather than the milliseconds of a mechanical seek, so even one outstanding request covers the 1000 images/s that the HDD's 100 IOPS cannot. Alternatively, converting the dataset to a sequential format (TFRecord/WebDataset) removes the random reads altogether and turns the HDD problem into a bandwidth one, in which the job needs 116 MB/s, within the roughly 150 MB/s that one HDD streams sequentially.

The first format choice is the record layout. Row-oriented formats, such as CSV or JSON lines, store successive records together and are easy to inspect and debug. They can stream many records from one file, although parsing text and reading unused fields consume CPU cycles and bandwidth.

A second choice solves a different storage problem. Columnar formats like Apache Parquet and Apache Arrow avoid unnecessary reads when tabular features can be selected before transfer. They organize data by column rather than by row, which allows for excellent compression as data within a column is typically of the same type and has lower entropy. Column pruning lets the reader fetch only required features. Predicate pushdown applies filters at the storage reader and can skip groups of records using their metadata, further reducing unnecessary I/O. These formats are well-suited for tabular ML tasks and feature engineering, such as the feature tables of Archetype B (DLRM at Scale), but are less natural for unstructured data like images or audio, which are treated as atomic blobs.

A third choice addresses the training access pattern directly. Sequential streaming formats reduce per-sample metadata to the cost of opening a shard. TFRecord, Mosaic's StreamingDataset format (Databricks 2026), and WebDataset[5] group many samples into large, contiguous binary files called shards. WebDataset uses standard `tar` archives, in which each sample keeps its image, label, and metadata entries while the loader streams the archive in order. A data loader opens a single shard and streams its contents sequentially, amortizing the open over thousands of samples. The `tar`-based formats have the added advantage of being valid targets for HTTP byte-range requests, making them ideal for streaming data directly from object storage without downloading the entire file first.

A key trade-off within streaming formats is the presence of an index. TFRecord readers often rely on a separate index file that records each sample's byte offset within a shard, while the scientific array format HDF5 embeds its index in the file itself. This enables efficient random access within a shard but introduces a metadata dependency, because the index must be read and held in memory. WebDataset is indexless, because each record in the `tar` archive contains its own header specifying its size, allowing for purely sequential streaming without any external metadata. The trade-off is

5 **WebDataset**: A data format that repurposes standard POSIX tar archives for ML training. The design insight is that tar's sequential header-then-data layout, invented in 1979 for tape backup, maps perfectly onto the streaming access pattern that ML data loaders need.

that indexless formats cannot efficiently seek to an arbitrary sample within a shard, a capability that matters for curriculum learning or active learning workloads that require nonsequential data access.

Compression integration also varies across formats, and it trades storage bytes against CPU decoding work. Parquet integrates compression at the column level, exploiting type homogeneity for high ratios. WebDataset and TFRecord typically leave compression to the individual sample. A WebDataset of images stores already-compressed JPEG files, while a text dataset might apply zstd compression per shard, and the host CPU pays for that decompression.

Consider two illustrative codecs on the same corpus: a lighter codec compresses 3:1 and produces 1.5 GB/s of decoded output per CPU core, while a denser codec compresses 4:1 and produces 0.5 GB/s per core. These are scenario assumptions, not fixed properties of named codec levels. At the 28 GB/s local NVMe rate, multiplying that read rate by each compression ratio yields potential decoded output rates of 84 GB/s and 112 GB/s, respectively, provided CPU decoding keeps pace. Dividing each uncompressed output rate by its per-core decoding speed gives 56 cores for the lighter codec and 224 for the denser one. The denser codec needs 4× as many cores, because it produces one-third more uncompressed bytes from the same storage read while decoding each byte at one-third the per-core speed.

Running each codec at the drives' full output is the worst case, an upper bound on decode work; a node needs only enough cores to decode its own share of demand. The generic text job of notebook 4.1 asks the entire cluster for about 167.8 MB/s, which a single core of either codec covers, so these core counts become the limit only for image and video feeds whose decoded demand per node approaches the drives' rate. With fewer cores than that demand requires, decoding limits delivery before storage saturates, so compression is chosen from the measured corpus, decoder, and available CPU budget rather than from the compression ratio alone.

For the running example, the 1.5 trillion token dataset is stored as roughly 2,000 shards of 3 GB each in a custom binary format. Each shard contains a small header (token count, vocabulary size, byte order) followed by a contiguous array of 4-byte token IDs. The number of `open()` calls per epoch is therefore the number of shards, not the number of tokens (trillions), an approximately 750,000,000× reduction in metadata operations that makes streaming from both local NVMe and remote storage feasible at training scale. When staged on storage with the standard Unix (POSIX) file interface, such as local NVMe or a parallel file system, the data loader memory-maps the shard, treating the file as a region of virtual memory, and indexes directly into the token array by position, achieving near-peak read bandwidth with zero per-sample metadata overhead. When the same shards are read directly from object storage, the loader uses HTTP ranged GET streaming and an explicit local buffer rather than memory mapping. The formats differ chiefly in the access cost each one imposes per sample on the storage path (table 4.4).

Table 4.4: Data Format Comparison for ML Training: Per-sample overhead determines whether the storage hardware's bandwidth is realized or wasted on metadata operations.

Format	Overhead/sample	Random access	Streaming	Compression
Individual files	≈ 50 µs (open/stat/close)	Yes	Poor	Per-file
Parquet	≈ 1 µs (row group seek)	Yes (row group)	Good	Column-level
TFRecord	≈ 0.1 µs (index lookup)	With index	Excellent	Per-record
WebDataset (`tar`)	~0 (sequential)	No	Excellent	Per-sample
Raw binary (tokens)	0	Yes (byte offset)	Excellent	None needed

Packaging removes the per-file cost within one node, but no node-local format gives every node the same view of the data. That requirement belongs to the next tier.

4.3.6 Tier 3 parallel file systems

Local NVMe drives cannot give 256 training nodes a single namespace for shared datasets or a globally visible home for checkpoints that survive a node crash. A parallel file system provides that shared tier across the cluster.

Definition 4.1: Parallel file system

Parallel file system (PFS) is a distributed storage architecture for ML training clusters that stripes data across many storage servers to provide aggregate throughput exceeding the capacity of any single device.

1. **Significance**: A PFS aggregates storage I/O bandwidth ($\text{BW}_{\text{storage}}$) across object storage servers (OSS). Twenty OSS nodes delivering 10 GB/s each offer 200 GB/s aggregate. When the client and network sustain that full rate, a 10 GB striped shard loads in about 50 ms rather than the 1 second a single 10 GB/s server would need. The rate that enters the data-movement term D_{vol}/BW of the single-machine iron law recalled in section 1.5.4 is the bandwidth that reaches the actual consumer, not the aggregate rating.
2. **Distinction**: Unlike network-attached storage, where every I/O request routes through a single server, a Lustre-style PFS client receives stripe location metadata from a dedicated metadata server (MDS) and then reads data directly from multiple OSS nodes in parallel. The MDS and OSS paths are architecturally separated, so data bandwidth scales with OSS count while metadata operations scale with MDS count.
3. **Common pitfall**: A frequent misconception is that a PFS has unlimited throughput if enough OSS nodes are added. In reality, every open, stat, and create on a Lustre-style PFS passes through the MDS, whose service rate is finite and does not grow with the number of OSS nodes; a one-time burst of 10,000 opens does not by itself establish sustained saturation. The MDS becomes a throughput bottleneck when the aggregate metadata-operation arrival rate exceeds its service rate, regardless of spare bandwidth at the OSS nodes.

A Lustre-style parallel file system decouples metadata lookups from bulk data movement (figure 4.2), separating two concerns that traditional file systems handle together (Schwan 2003). Metadata servers (MDS) manage the namespace (persisting file layouts and attributes on metadata targets, or MDTs): file creation, directory listings, permission checks, and lock management. Object storage servers (OSS)[6] manage the actual data blocks (stored on object storage targets, or OSTs), each serving a stripe of every large file. When a client opens a 10 GB training shard, the MDS tells the client which OSS nodes hold which stripes, and the client reads from all of them in parallel. A Lustre[7] deployment with 100 OSS nodes, each providing 10 GB/s, delivers an aggregate 1 TB/s.

This abstract architecture has concrete implications for performance tuning. When a training job creates a new dataset directory on Lustre, the administrator configures the stripe count (the number of OSTs a file is spread across) and stripe size (the chunk size written to each OST) for that directory (Lustre Wiki 2017). For large sequential-read training shards, administrators may choose wider striping and multi-megabyte stripe sizes so a single file can draw bandwidth from multiple OSS nodes. Lustre's striping guide treats wide striping as useful for very large files or files accessed by many clients, while smaller files often use fewer stripes to reduce overhead (Lustre Wiki 2019).

A 4 GB data shard striped across 100 OSTs with 4 MB stripes places roughly 10 stripes on each OST. When a data loader reads large contiguous ranges, the Lustre client can issue parallel read requests to the OSS nodes holding those stripes, aggregating their bandwidth. The client also maintains its own read-ahead buffer, prefetching the next several stripes while the application processes the current ones. This file-system-level read-ahead is distinct from the data loader's application-level prefetch buffer; the two layers of prefetching compound and can hide much of the delay to the OSS nodes from the training process, provided delivery keeps pace.

The separation between metadata and data paths is how a Lustre-style system reaches the aggregate bandwidth that ML workloads demand. GPFS[8] (IBM Spectrum Scale) is a shared-disk alternative that stripes data and metadata across its disks and coordinates concurrent access through token-based distributed locking (Schmuck and Haskin 2002). Its lock protocol preserves consistency across clients without a simple MDS/OSS split, at the price of more complex lock management that must scale to thousands of nodes.

The data path scales with the number of object storage servers, but every file still begins with a request to the metadata service, and that service has a finite operation rate. That finite rate is

6 **OSS (object storage server)**: In parallel file system terminology, "object" means a chunk of striped file data managed by an object storage target, a usage that predates cloud object storage (S3, GCS) by over a decade. Confusing the two is a common source of miscommunication: when a Lustre administrator says "add more OSS nodes," they mean adding data-serving capacity to the parallel file system, not provisioning cloud buckets.

7 **Lustre**: A portmanteau of "Linux" and "cluster," developed at Carnegie Mellon University and first deployed in production in 2003. Lustre is common in HPC and ML infrastructure because its architecture scales aggregate bandwidth linearly with the number of OSS nodes, the same property that makes it a natural choice for training clusters where a single job may demand hundreds of GB/s of sustained read throughput.

8 **GPFS (General Parallel File System)**: Developed by IBM Research starting in 1998, now marketed as IBM Spectrum Scale. For ML checkpoint writes, its token-based distributed coordination can reduce reliance on a single dedicated metadata path, but it also makes lock behavior part of the performance envelope.

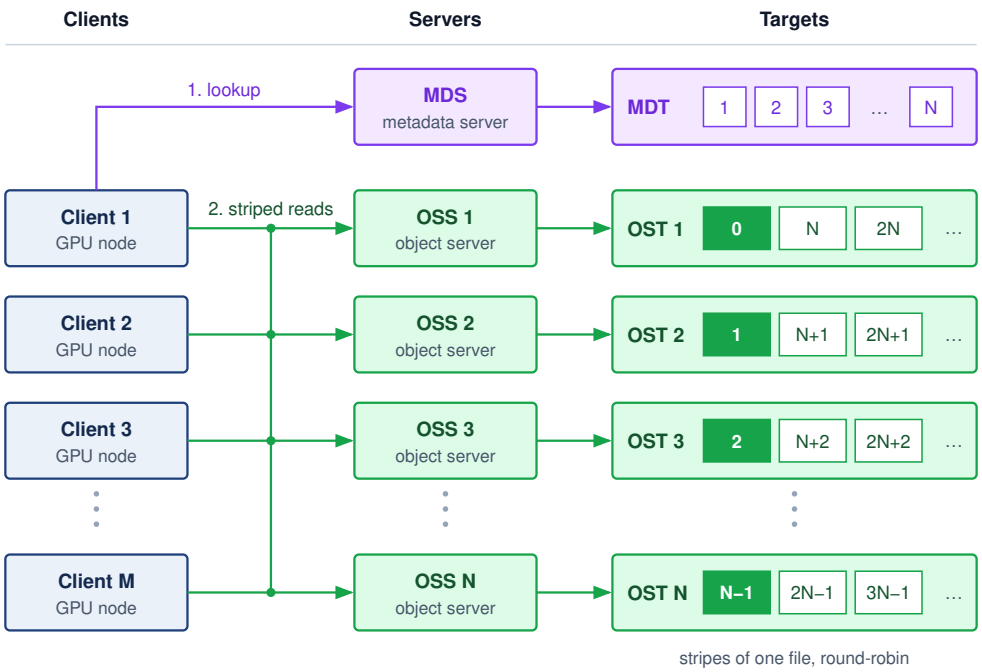

Figure 4.2: Parallel File System Architecture: A client makes one metadata request (MDS → MDT) to learn a file's stripe map, then reads the stripes directly from the object storage servers (OSS → OST) in parallel. Each OST holds every Nth stripe of the file in round-robin order, so consecutive stripes (dark) sit on different servers and one client reads them at once. The shared data path lets every client draw on every OSS, so aggregate bandwidth grows with the number of servers while metadata traffic stays on a separate path.

the source of the small file problem. The per-open cost that section 4.3.5 charged to one node's CPU now lands on a metadata service shared by the whole fleet. Each `open()` system call requires a metadata lookup, lock acquisition, and attribute fetch. Suppose a dataset of millions of 10 KB images is stored as individual files, and each of 10,000 workers, a fleet several times the running cluster, opens a new file every few samples. The aggregate open rate then exceeds the few hundred thousand operations per second an MDS typically serves. The MDS becomes the serialization point, and delivery collapses to its operation rate long before the data links saturate.

Definition 4.2: Small file problem

Small file problem is an ML data-loading pathology in which per-file operations overwhelm a storage system's metadata service.

1. **Significance**: `open`, `stat`, and `close` requests from 10,000 workers can serialize at the metadata service, idle the cluster, and reduce effective I/O bandwidth ($BW_{storage}$).
2. **Distinction**: This is a metadata request-rate and latency (L_{lat}) problem, not a bulk bit-rate problem.
3. **Common pitfall**: Faster SSDs do not fix this format problem. Packing samples into sequential containers such as TFRecord or Parquet amortizes metadata operations across many samples.

At production scale, the metadata bottleneck can collapse an otherwise high-bandwidth file system. When a cluster migrates from aggregated containers back to individual sample files, metadata request queues saturate long before the underlying storage servers or network links reach their byte-rate limits.

Example 4.1: The metadata meltdown

Scenario: A large-scale training cluster migrates from a preprocessed sequential dataset to 200 million raw image files stored individually on a parallel file system provisioned for 500 GB/s aggregate bandwidth.

Diagnosis: The file system delivers less than 1 percent of rated throughput. Millions of per-second `stat()` and `open()` calls overwhelm the metadata servers, choking data loaders despite abundant network and storage bandwidth.

Systems lesson: Metadata contention limits delivery when operation demand exceeds metadata service capacity. Bundling 200 million raw images into 50,000 large tar shards (for example, WebDataset) reduces the file count by 4,000× and amortizes per-file operations across samples. Packaging alone does not prevent startup contention if loaders still enumerate files with per-file `stat()` calls; distributing a precomputed dataset manifest removes those repeated lookups from the hot path. Measure metadata operation rate alongside byte bandwidth to determine which service limits delivery.

Striping and aggregation solve different problems. Striping distributes large files across data servers, whereas aggregation, the shard packaging of section 4.3.5, reduces the number of files the MDS must track. Realized bandwidth then depends on the reader. A read contained within one stripe uses a single server, while enough concurrent reads spanning stripes can draw from several servers at once. Wider striping therefore helps only when read sizes and concurrency exploit those paths, and small files usually need fewer stripes.

Parallel file systems provide a shared namespace and lock management that support coordinated checkpoints (Schmuck and Haskin 2002), but the two ML write patterns need different guarantees from them. Immutable training shards remove conflicts with concurrent writers once the dataset version has been published, so readers need only agree on which completed version belongs to the run. A checkpoint needs more. File-system locks coordinate individual writes, and a completed write to one file does not establish that the other shards are complete, so a multi-file generation needs an application-level publication protocol, which declares it restorable only after every shard has persisted. Section 4.6 states that protocol.

The performance of a parallel file system depends on the *balance* of load across OSS nodes as well as their *number*. If a shard uses 10 of 100 equally provisioned OSS nodes, those servers offer at most 10 percent of the system's aggregate server bandwidth. Realized delivery can be lower because of client limits, request size, or insufficient concurrency. Conversely, when many concurrent readers draw on shards striped across all 100 servers, demand spreads over every server and realized bandwidth can approach the system's rating, provided the clients and network keep pace. Placement and access

scheduling therefore matter alongside the number of storage servers, and a simple shard-assignment calculation shows how readers can bunch on one shard even when aggregate storage capacity looks ample.

↪ Tail latency as a serving constraint, where routing and load balancing decide which requests miss the p99 target, is developed in section 10.11.

Napkin Math 4.4: The thundering herd: Shard contention

Problem: A dataset is split into 1,000 shards on a shared file system. If 32 workers each pick a shard at random to start their next epoch, what is the probability that at least two GPUs "collide" on the same shard, creating a potential read hotspot? A shard collision does not imply a server collision unless the placement maps each shard to a distinct server.

Math: This is a variant of the "Birthday Problem" in probability.

1. **Probability of No Collision**: $\approx e^{-n_{\text{workers}}^2/(2K_{\text{shards}})} = e^{-32^2/2{,}000} \approx 0.60$.
2. **Probability of Contention**: $1 - 0.60 = 40\%$.

Systems insight: With 32 workers and 1,000 shards, the 40.1 percent collision probability identifies a potential storage-shard hotspot. Whether that collision delays workers depends on shard placement, caching, and available service capacity. Coordinating shard assignments can reduce this source of tail latency when assignments spread demand across independently serving devices while preserving the intended sampling policy.

At the scale of thousands of nodes, tail latency (Dean and Barroso 2013a) can dominate exposed input wait. If a training step requires data from 100 storage servers and each independently has a 1 percent chance of being slow, the probability that all 100 respond quickly is $0.99^{100} = 0.366$, meaning over 63 percent of these request sets contain at least one slow response. Shared congestion can correlate delays, changing that probability. A slow response delays the step only when prefetching fails to hide it and the required input lies on the step's critical path. When it does, the delay is not confined to one worker, because synchronous data-parallel training waits for every worker at the next collective, so one late fetch stretches the whole fleet's step.

Systems can mitigate tail latency through hedged requests: after a chosen delay, the client issues a redundant read to another replica and uses the first valid response. The trigger delay trades faster tail completion against extra work (Dean and Barroso 2013b). An independence model can multiply the probabilities of two slow replicas, but shared congestion and the later start of the hedge invalidate a simple squared-probability latency guarantee. Issuing a hedge at the median response time would duplicate roughly half of requests before accounting for cancellation; that load can worsen the congestion the scheme aims to avoid. A hedge therefore pays off only when its trigger sits in the measured tail and the duplicate traffic it adds stays within the spare service capacity.

The effectiveness of hedged requests depends on the replicated data layout. If both replicas reside on OSS nodes that share the same network switch, a switch failure makes both copies simultaneously unavailable. Effective hedging requires failure-domain-aware placement: replicas should reside in different racks, connected through different top-of-rack switches, so that the failure of any single component affects at most one replica.

In production clusters, the parallel file system is a shared utility, simultaneously servicing many concurrent training jobs with different I/O patterns. At any given moment, a large language model job might be streaming sequentially through text shards, a computer vision job might be reading random image shards, and a checkpoint storm from a third job could be saturating write bandwidth. This concurrency gives rise to the noisy neighbor problem, where one job's I/O pattern severely degrades performance for all other jobs. A checkpoint save that consumes a disproportionate share of the PFS's internal network bandwidth causes read latency for other jobs to spike, potentially stalling their training pipelines. The impact of sharing is easy to quantify. A PFS with 1 TB/s aggregate bandwidth, shared across 10 concurrent training jobs, provides only 100 GB/s per job on average. If one job's checkpoint storm consumes 400 GB/s for 10 seconds, the remaining nine jobs share 600 GB/s, a 33 percent reduction that can trigger data stalls in their accelerator pipelines. PFS administrators mitigate this through I/O scheduling policies (quality-of-service tiers, bandwidth quotas per job), but these policies add administrative complexity and can reduce peak single-job throughput.

A related challenge is namespace isolation, which gives each team its own directory tree with its own striping settings. Different teams and workloads require different storage configurations. A team training on millions of 100 KB images first needs to aggregate samples to reduce metadata

operations; increasing stripe count does not remove per-file metadata work. A team working with large video shards can use wider striping when read sizes and concurrency exploit the additional data paths. Misconfigured striping for one team's workload can create hotspots that degrade performance for the entire shared file system.

Checkpoint 4.2: Parallel file system design

Consider a training cluster with 512 nodes, each running 8 GPUs. The training job requires 400 GB/s of aggregate read bandwidth.

- ☐ **OSS count**: calculate the number of OSS nodes required when each Lustre OSS delivers 10 GB/s.
- ☐ **Metadata cost**: calculate the time required to open 200 million 116 KB image files sequentially when each `open()` call takes 50 µs on the MDS, then state the practical implication.
- ☐ **Shard count**: calculate the number of metadata operations needed to read the full dataset when images are bundled into 4 GB tar files.

Parallel file systems solve the shared-namespace problem for active training clusters, but their cost per byte is too high for the organization's complete data holdings. The petabytes of raw training data, historical checkpoints, and preprocessed datasets that an ML organization accumulates over years require a tier optimized for capacity and durability rather than bandwidth.

4.3.7 Tier 4 object storage

Keeping petabytes of raw corpora, preprocessed shards, and historical checkpoints on a parallel file system would pay high-bandwidth server rates for data that sits mostly at rest. At petabyte scale, object storage provides the durable, lower-cost foundation of the hierarchy. Services like Amazon S3 and Google Cloud Storage organize data in object namespaces where objects are retrieved by keys rather than by POSIX directory paths (Amazon Web Services 2026c; Google Cloud 2026a). Unlike file systems, which organize data in hierarchical directories with metadata-rich operations (rename, link, permission inheritance), object storage treats each object as an opaque blob identified by a unique key. The simpler interface changes the namespace operations exposed to clients; the service still manages metadata for placement, versions, access control, and integrity as capacity grows.

Services such as Amazon S3 publish an "eleven nines"[9] annual durability target (99.999999999 percent). Erasure coding,[10] replication, placement, and repair contribute to storage protection, but a coding scheme alone does not establish a service's durability target. Selecting a source-of-truth tier for training data and checkpoint archives requires checking the chosen service's failure model alongside versioning, lifecycle policy, access control, and recovery testing.

Three different guarantees govern whether that source of truth is usable. Durability asks whether acknowledged data survives failures; availability asks whether the storage service responds when a client needs it; and consistency asks which version a read is allowed to observe when writes overlap or replicas disagree. Replication can improve durability and availability without by itself guaranteeing that a reader sees the newest committed checkpoint. A recovery design must therefore state all three properties instead of treating a durability percentage as a complete reliability contract.

9 **Eleven nines (99.999999999 percent durability)**: A design target of 11 nines annual durability corresponds to an expected loss probability of about 10^{-11} per object-year; for one million objects, that is about one expected object loss per 100,000 years.

10 **Erasure coding**: Built on Reed-Solomon codes (Reed and Solomon 1960), which use maximum-distance-separable coding to recover the original data from any sufficient subset of encoded fragments. In storage systems, Reed-Solomon-style codes and later local-reconstruction variants trade storage overhead against repair bandwidth and degraded-read cost (Plank 1997; Huang et al. 2012). Local-reconstruction codes, for example, read fewer surviving fragments to rebuild a lost one, cutting repair traffic at the price of a little extra capacity.

Definition 4.3: Erasure coding

Erasure coding protects stored data by adding redundant fragments, with the Reed-Solomon scheme considered here encoding d data and m parity blocks so that any d remaining fragments recover the original.

1. **Significance**: Coding can protect against fragment loss with less capacity overhead than full replication. A 4+2 Reed-Solomon code stores four data fragments and two parity fragments, tolerating two unavailable fragments at 1.5× storage overhead instead of the 3× overhead of triple replication. Actual durability also depends on failure correlation, placement, detection, and repair.

2. **Distinction**: Unlike replication (which stores complete copies), erasure coding uses mathematical encoding (for example, Reed-Solomon) to distribute redundant information across different failure domains (disks, racks, or sites).
3. **Common pitfall**: A frequent misconception is that erasure coding is "free durability." In reality, it is a latency-compute trade-off: reconstructing data after a failure requires additional CPU cycles and increases the p99 latency of the storage read.

Erasure coding lets the service recover data without requiring the training application to manage fragments. In the form considered here, the service retains the original d data fragments alongside parity, so healthy reads can use the data fragments directly. When a required fragment is unavailable, reconstruction uses surviving data and parity fragments. The object interface hides that recovery procedure, but additional I/O, network traffic, and decoding can remain visible as longer read latency. Protection against data loss does not imply unchanged service time during a failure.

Object storage also costs less per byte. At a representative $0.02/GB/month, its recurring capacity rate is 5× lower than the local-NVMe rate of table 4.2, and section 4.8 converts that ratio into annual costs for a 100 TB dataset. Comparing it with HBM acquisition cost additionally requires a hardware lifetime and utilization assumption. This cost advantage makes object storage a common home for the organization's training data lake, where raw data, preprocessed datasets, and archived model artifacts can reside durably.

Request latency also matters. Design estimates use 50 to 100 ms for an object GET as a representative planning range rather than a universal service guarantee. Much of that delay is the price of the indirection that provides durability. The service looks up the object's location in a distributed metadata index, determines which fragments to read, fetches them from potentially different storage nodes, and reconstructs data when a fragment is unavailable, all on top of the client network path and queueing. Large transfers amortize per-request overhead across more bytes, while small transfers can spend most of their elapsed time on that overhead. Request size and concurrency therefore enter the delivery budget alongside aggregate bandwidth.

The remedy is the shard packaging of section 4.3.5, applied here to request latency rather than metadata operations. Formats like WebDataset (Aizman et al. 2019) or Mosaic Streaming (Databricks 2026) bundle samples into multi-gigabyte shards that transform high-latency random access into high-bandwidth sequential streaming. Rather than issuing billions of small GET requests, the data loader issues hundreds of large reads in parallel, each fetching an entire shard that contains thousands of samples. With sufficient parallelism and colocated network capacity, object storage can sustain high aggregate throughput, enough to feed large training clusters for some workloads. The streaming pattern works by assigning each data loader worker a subset of shards and having it read them sequentially. Within each shard, samples are stored contiguously, so the worker decodes them in order and shuffles them in a local buffer. This approach achieves pseudo-random access across the dataset while maintaining sequential I/O at the storage level.

The architecture of a streaming data loader for object storage differs from a local-storage loader in important ways. A local-storage loader can use memory-mapped I/O to treat the dataset as a virtual memory region, relying on the operating system's page fault mechanism to load data on demand. This approach is elegant but incompatible with object storage, which answers HTTP requests rather than offering the POSIX file system interface that memory mapping requires. Instead, an object-storage loader must explicitly manage HTTP connections, issue ranged GET requests for specific byte ranges within shards, handle retries on transient failures, and manage a local buffer pool for decoded samples. Libraries like WebDataset and Mosaic Streaming encapsulate this complexity, presenting a simple iterator interface to the training loop while managing the HTTP transport and buffering internally.

Applying this streaming architecture to the running example makes the design concrete as an alternative to the staged path the chapter adopts. Instead of copying shards to local NVMe at job start, the cluster would read them from the object store throughout the run. The tokenized training set is stored in an S3-like object store as roughly 2,000 shards of 3 GB each. The streaming architecture assigns each of the 256 compute nodes a unique subset of roughly 8 shards. Within each node, 8 data loader worker processes issue concurrent HTTP `Range` requests to fetch 256 MB chunks from their

assigned shards. Across the cluster, this results in 2,048 concurrent requests, allowing aggregate throughput in the 50–100 GB/s range when the object store and network path are provisioned for it, far exceeding the roughly 167.8 MB/s required by the generic text job in the text-versus-image comparison. The running example's average input demand is lower still. The network between the compute cluster's virtual private cloud (VPC, its isolated network inside the provider) and the object-storage endpoint can cap streaming throughput, but at these provisioned service rates the cap still exceeds what the running job needs. Organizations that co-locate compute and storage in the same cloud availability zone (one isolated data-center site of a provider) can reduce cross-network latency, which lowers the required prefetch buffer depth for jobs with short steps and improves overall pipeline efficiency.

That request parallelism is what object storage rewards. Each of the 2,048 GET requests is independent and can be served by a different storage node within the cloud provider's infrastructure, and the count grows with the fleet. With each request fetching a 256 MB chunk of a shard, the aggregate throughput depends on the cloud provider's network capacity, internal bandwidth, and request-rate quotas. The practical limit is often the network bandwidth between the compute cluster and the object storage service, not the storage service's internal media bandwidth.

Many object storage services provide strong read-after-write consistency, which means that a subsequent read of an acknowledged object sees the completed write. Amazon S3 announced strong read-after-write consistency for object writes and overwrites, together with strongly consistent listing, in December 2020 (Amazon Web Services 2020). The relevant contract is the documented behavior of the chosen backend, not a blanket assumption about object storage. Per-object consistency still does not make a checkpoint spanning many objects atomic, so the publication protocol of section 4.6 applies here as it does on a file system.

The scale of ML datasets makes them vulnerable to silent data corruption: subtle bit flips in the storage medium or during network transfer that can introduce training artifacts nearly impossible to diagnose. Object storage services provide per-object checksums that verify integrity on upload, copy, or download; Amazon S3, for example, documents checksum options including CRC-family and cryptographic hashes (Amazon Web Services 2026a). Parallel file systems often rely on underlying RAID or erasure coding for physical integrity but may not provide end-to-end checksums visible to the application. A well-engineered large-scale ML pipeline computes and stores a checksum for each data shard during preprocessing and verifies it when the shard is first loaded by a training worker. The verification overhead is negligible compared to the cost of training on corrupted data, and checkpoints need the same check, because a flipped bit in restored state can make training diverge silently after recovery.

↪ Detecting and verifying silent data corruption across the fleet, including in restored checkpoints, is developed in section 7.5.

Training datasets can contain sensitive information, including personal data, proprietary text, and licensed images. Storage design must carry the organization's access and retention policies through every copy. Object services expose identity policies and encryption options; parallel file systems provide POSIX permissions and access-control lists. Audit coverage and enforcement depend on the configured platform, not the tier's name. Versioned manifests should identify derived shards and retained copies so that access reviews, retention decisions, and deletion workflows can locate the relevant artifacts; a deletion request, for example, must reach every immutable preprocessed shard derived from the deleted record. Storage controls protect stored bytes; evaluating what a trained model retains requires additional model-level analysis.

The transition from Tier 4 to Tier 5 is driven primarily by economics. Data that is accessed rarely should move to archive storage when retrieval latency and governance requirements allow it, because archive capacity is cheaper still.

4.3.8 Tier 5 archive and cold storage

Archive services such as S3 Glacier preserve rarely accessed lineage (superseded checkpoints, raw source corpora, and training logs) over multi-year horizons so that a deployed model can still be reproduced or audited years after its run ends. At a representative $0.004/GB/month, archive storage is 5× cheaper than the $0.02/GB/month object-storage assumption used in section 4.3.7, trading that lower capacity rate for retrieval latency measured in minutes to hours and making this tier unsuitable for operational access. The record at stake is training data provenance: the link from a deployed model to the exact data versions, hyperparameters, and preprocessing that produced it.

The organization's governance policy determines which records to retain and for how long, and the storage design must keep them identifiable and retrievable for that period; archive serves when its retrieval delay meets the requirement, while records needed during frequent investigations may justify a warmer tier.

Automated lifecycle policies govern the transition of data between tiers based on access recency. As one illustration of how such automation is expressed, a time-based rule for ML infrastructure might move checkpoints as follows. The most recent three checkpoints remain on the parallel file system for fast recovery. Checkpoints older than 72 hours are transitioned to object storage, and checkpoints older than 90 days are transitioned to archive storage. Checkpoints older than two years are deleted unless flagged for regulatory retention. S3 Lifecycle rules provide the same general transition/archive/delete mechanism for object lifetimes (Amazon Web Services 2026b). The running example's checkpoint retention in section 4.6 adopts a count-based variant of the same idea, where the transition rule keys on checkpoint count rather than age. Without automation, teams either hoard data on expensive tiers (wasting budget) or delete data too aggressively (losing reproducibility). The lifecycle policy encodes the organization's cost-recovery trade-off in a declarative rule that requires no manual intervention.

For the 175B model training run, the archive tier holds the complete training lineage: the raw 3 TB compressed source corpus, the 6 TB serialized shard set, all preprocessing scripts, the final checkpoint, and a sampled subset of intermediate checkpoints. Capacity must include every retained full checkpoint as well as both corpus representations; table 4.5 computes that total for the declared retention policy. The asymmetry between archive cost and compute cost makes retention inexpensive relative to the run it documents; failures to archive usually reflect missing lifecycle ownership, unclear retention policy, or retrieval-process neglect rather than raw capacity cost.

Archive storage services offer multiple retrieval speed tiers, each with different pricing. In an S3 Glacier-style service, standard retrieval may complete within hours at moderate cost, expedited retrieval may complete within minutes at a higher cost, and bulk retrieval may trade still more latency for the lowest retrieval price (Amazon Web Services 2026d). The choice of retrieval tier depends on the urgency of the use case. Forensic analysis of a production incident warrants faster retrieval; annual compliance audits can use slower bulk retrieval. Designing the lifecycle policy requires understanding the storage cost, the expected retrieval frequency, and the urgency for each class of data.

The six tiers now have their roles, capacities, and prices, but a role does not say how fast bytes must climb the hierarchy. That rate is set by the accelerators consuming them, and it is the next quantity to derive.

4.4 The Data Pipeline Equation

Whether a training job streams from object storage or reads from local NVMe, each worker's next batch must reach HBM before the current compute step finishes; falling short starves the fleet's accelerators, while provisioning far above demand strands capital in idle storage bandwidth. Sizing the hierarchy therefore begins by turning the accelerator consumption rate into a required byte rate. The data pipeline throughput equation in equation 4.2 quantifies the bandwidth the pipeline must sustain for a target training rate:

$$\mathrm{BW}_{\mathrm{required}} = N_{\mathrm{GPU}} \times \eta_{\mathrm{target}} \times \frac{D_{\mathrm{vol}}^{(\mathrm{batch})}}{T_{\mathrm{compute}}} \tag{4.2}$$

where N_{GPU} counts accelerators consuming distinct local input batches, η_{target} is the target fraction of wall-clock time spent computing (typically 0.8 to 0.95), $D_{\mathrm{vol}}^{(\mathrm{batch})}$ is one local batch in bytes, and T_{compute} is its uninterrupted compute time. Devices that cooperate on one batch do not each require an independent storage read, so their shared input must be counted once at the tier being sized.

The target fraction decides which rate the pipeline must meet. Eliminating sustained input stalls requires sizing the pipeline to full demand, which supplies batches at the full compute rate. A target fraction below one sizes it for a slower training rate that the team has chosen to accept. This target is an active-time fraction, not model FLOPs utilization. If iteration time already includes stalls, multiplying by the target fraction again would count the slowdown twice. The text-versus-image

comparison of notebook 4.1 already applied this equation with $\eta_{\text{target}} = 1$; the equation makes the target fraction explicit, and the service chain below adds the requirement that every stage, rather than the accelerator-facing endpoint alone, meet the resulting rate.

The demand equation sizes the accelerator-facing endpoint, but bytes must traverse the complete service chain. Let r_j denote the measured service rate of stage j, including object-store reads, network transfer, decoding and augmentation, host-memory copies, and host-to-device transfer. Express all rates in the same output-equivalent bytes per second, using each stage's expansion or compression ratio; raw compressed-read bandwidth cannot be compared directly with decoded-tensor bandwidth. With bounded buffers and no sustained backlog, the pipeline can deliver at most

$$r_{\text{pipeline}} = \min_j r_j, \qquad r_{\text{pipeline}} \geq \text{BW}_{\text{required}}. \tag{4.3}$$

Capacity at one tier cannot compensate for a slower downstream stage. Compression may raise the storage-read rate while lowering the decode rate; additional loader workers may raise decoding throughput while saturating the host-memory or PCIe stage. GPUDirect Storage (section 4.5) can drop the host-copy stage from the chain, which raises the minimum when that copy is the limiting stage or takes cores from the stage that is. The end-to-end inequality in equation 4.3 therefore turns the chapter's hierarchy into a testable contract. Measure every stage under representative concurrency, then provision or overlap the stage whose service rate is the minimum.

The throughput equation in equation 4.2 exposes four levers on storage demand, and none of them is free. Reducing the accelerator count (N_{GPU}) lowers aggregate input demand by reducing the number of distinct consumers, but it lowers training throughput with it, whereas lowering target utilization (η_{target}) relaxes the delivery rate only by accepting idle accelerators. Decreasing the per-batch data volume ($D_{\text{vol}}^{(\text{batch})}$) through a more compact representation of the same samples helps only while decoding keeps pace, and shrinking the batch itself changes the training configuration and its statistical efficiency. Increasing compute time (T_{compute}) gives storage longer to deliver each batch, at the cost of proportionally longer iterations. This is also why a faster accelerator can expose an input bottleneck that a slower one hid. The equation therefore sizes the pipeline for a chosen training configuration; deliberately slowing useful computation does not improve delivery.

The consequences of falling short are captured by the **data stall ratio**, defined in equation 4.4 as the fraction of each training step where the accelerator waits for data:

$$\text{Data Stall \%} = \frac{T_{\text{step}} - T_{\text{compute}}}{T_{\text{step}}} \times 100 \tag{4.4}$$

The data stall ratio covers one worker's input path with communication excluded, so its T_{step} is compute plus exposed input wait. The Fleet Law of section 1.5.4 carries no input-wait term, so it implicitly treats the next batch as already resident. Whatever input wait the loader fails to hide adds to the distributed step on top of $T_{\text{comm}}(N)$ and $T_{\text{sync}}(N)$, as compute (C_1) capacity that sits idle while storage catches up. In a real distributed job, the gap $T_{\text{step}} - T_{\text{compute}}$ also contains exposed communication and synchronization, so the ratio implicates storage only when that gap is measured as time spent waiting on the input queue, separately from collectives and barriers. Without overlap, step time adds compute and I/O time. Pipelining reduces step time to the longer of the two durations and stall time to the positive difference between I/O and compute.

Napkin Math 4.5: The data stall ratio

Problem: A GPU processes a batch in 200 ms, but storage takes 250 ms to deliver the next batch. What fraction of each training step is the accelerator idle, and how much does pipelining I/O with compute reduce that stall?

Without pipelining (sequential I/O then compute):

$$T_{\text{step}} = T_{\text{IO}} + T_{\text{compute}} = 250 + 200 = 450 \text{ ms}$$

$$\text{Stall \%} = \frac{250}{450} = 55.6\%$$

With pipelining (I/O overlapped with compute):

$$T_{\text{step}} = \max(200, 250) = 250 \text{ ms}$$

$$\text{Stall \%} = \frac{250 - 200}{250} = 20\%$$

Systems insight: Even with pipelining, the accelerator is idle 20 percent of the time. To eliminate the stall entirely, the next batch must arrive within the 200 ms compute window, either by increasing effective storage bandwidth or by adding parallel I/O streams. Deeper prefetch buffers cannot help here, because the delivery time exceeds the compute window on every step rather than merely in the tail.

The data stall ratio measures time lost waiting for inputs relative to the chosen iteration window. Acceptable stall time depends on the fleet's cost and throughput objectives; a threshold is a policy choice, not a universal storage-health boundary. Instrument batch-ready times and accelerator execution in the same trace to distinguish input starvation from synchronization, CPU preprocessing, and other idle intervals. General profiling tools can expose these events, but an idle accelerator alone does not identify storage as the cause.

A storage pipeline that met its budget at launch can drift out of it as datasets grow, jobs contend, and drives age, so it requires continuous monitoring to keep silent bottlenecks from eroding training efficiency. Key metrics to track in real time include per-tier bandwidth utilization (NVMe, network, object storage), prefetch buffer depth (how many batches are queued and ready), the data stall ratio per GPU, the parallel file system's metadata operation rate, and NVMe drive health indicators (wear level, temperature, error counts). Alert thresholds should trigger before stalls become visible in training loss curves. An undetected 5 percent increase in data stall ratio across the 1,024-GPU reference cluster (half the running cluster) for just one week wastes about $17,203 in idle compute (1,024 GPUs $\times$ $2/GPU-hour $\times$ 168 hours $\times$ 5 percent). Comprehensive storage monitoring is therefore a first-order economic concern rather than an operational nicety.

↪ An I/O-wait efficiency target for the fleet Data axis is developed in section B.10.4.

Diagnose a storage stall at three levels. The application timeline identifies when the accelerator waits; per-worker rates and queue depth identify the pipeline stage that falls behind; device and shared-server counters show which resource limits that stage. Correlating those observations distinguishes a byte-bandwidth shortage from metadata contention or insufficient decoding capacity. Because a small development dataset may hide the contention and accumulated state of a production run, a pipeline stress test should use representative data volume and concurrency and run long enough to expose queue depletion, resource leaks, and memory pressure before the loader is coupled to an expensive training job.

The terms in equation 4.2 explain why the chapter's workloads place such different demands on the same cluster. In the bandwidth comparison of notebook 4.1, image classification with ResNet-50 on ImageNet combines large compressed samples (116 KB per image on average) with short iteration times (100 to 200 ms), pushing $D_{\text{vol}}^{(\text{batch})}/T_{\text{compute}}$ high, whereas language-model pretraining packs 4,096 4-byte token IDs into 16 KB per sequence even though total volume over the run still reaches trillions of tokens. Recommendation model training, the Archetype B (DLRM at Scale) workload, combines both access modes. Embedding lookups require high-IOPS random access to the embedding table while the dense layers stream sequential training batches. Which storage tier limits delivery therefore follows the workload's access profile rather than cluster size alone. Archetype C (Federated MobileNet) falls outside this shared-storage analysis, because its raw training data stays on the devices that produce it; there is no central corpus to stage, and its storage constraints become per-device capacity and upload limits.

MLPerf Storage measures these profiles on production storage systems. Each submission emulates a set of accelerators that consume samples with a fixed per-sample compute time, so the storage system, not a real GPU, sets the pace. Each submission reports the read throughput the system sustained while keeping the emulated accelerators at least 90 percent busy (70 percent for CosmoFlow,

↪ On-device training data and federated updates are developed in chapter 11.

which trains on 3D cosmology simulation data) (MLCommons 2024; Devarajan et al. 2021). Across the 124 passing results of the closed division in the 2024 round, sustained throughput grows in proportion to the number of emulated accelerators, as the N_{GPU} term of equation 4.2 predicts (figure 4.3, panel a).

Per accelerator, the median demand differs about 15× across workloads. ResNet-50 is demanding by image-classification standards, yet it is the lightest of the three because its samples are small 2D images while 3D U-Net reads whole medical volumes. An emulated H100 needs about 2.8 GB/s for 3D U-Net on medical volumes, 0.55 GB/s for CosmoFlow, and 0.18 GB/s for ResNet-50 (figure 4.3, panel b). Emulating H100s instead of A100s raises every workload's demand by 1.6–2×, because the faster accelerator shortens T_{compute} while the bytes per batch stay the same. A storage tier sized for one accelerator generation is therefore undersized for the next, even when the dataset and the model do not change.

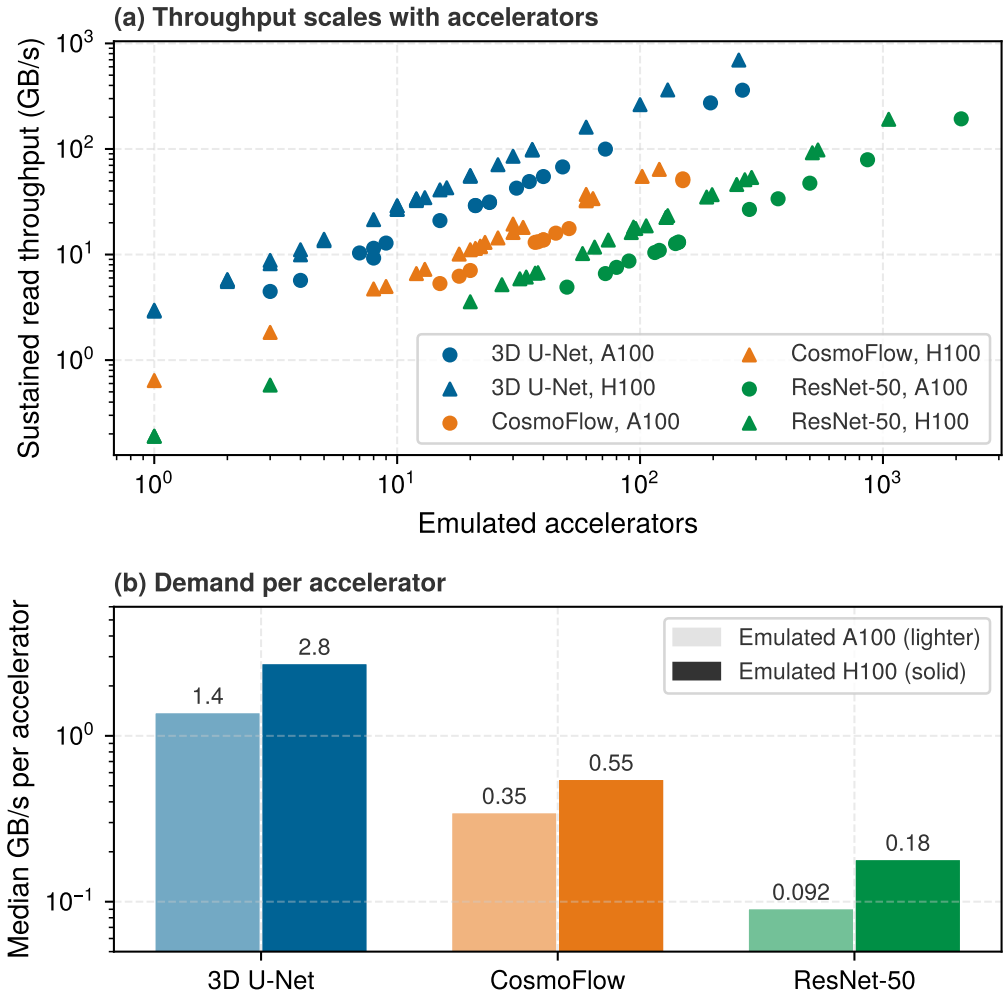

Figure 4.3: Measured Storage Demand per Accelerator: (a) Read throughput sustained by each passing closed-division MLPerf Storage v1.0 submission against its number of emulated accelerators, on logarithmic axes, by workload (color) and emulated accelerator (marker); throughput grows in proportion to the accelerator count. (b) Median throughput per emulated accelerator for each workload and accelerator, on a logarithmic axis. Accelerators are emulated with a fixed per-sample compute time and must stay at least 90 percent busy (70 percent for CosmoFlow); the values measure storage systems under that load, not real GPUs. Data from MLCommons (2024).

In an ideal overlapped pipeline, the bandwidth-to-demand ratio bounds the compute active-time fraction (figure 4.4). With input demand defined as batch bytes divided by uninterrupted compute time, the frontier ramps linearly through bandwidth-starved operating points before flattening at the supply-equals-demand saturation knee (ratio = 1). The active fraction is the smaller of one and available bandwidth divided by that demand, assuming no other bottleneck.

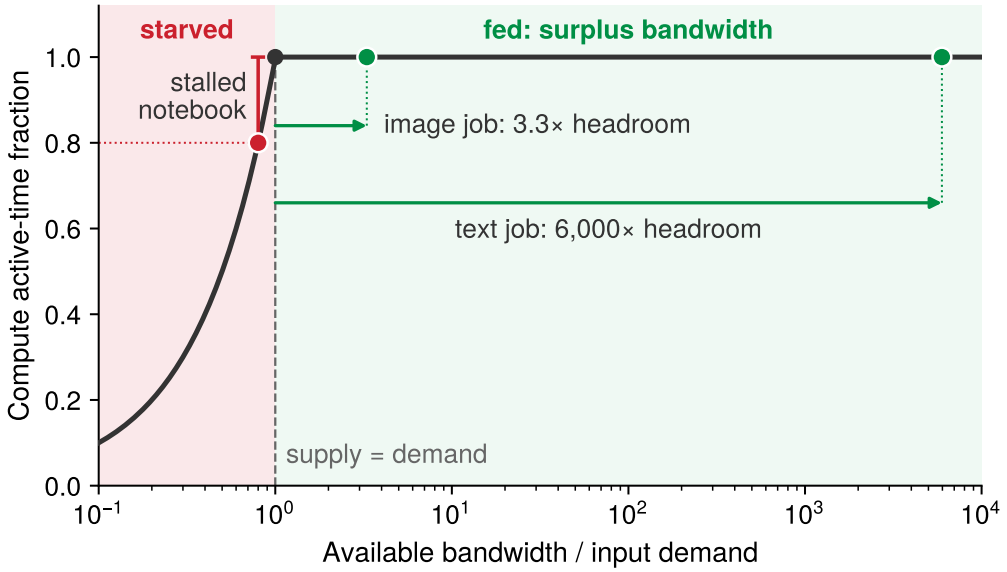

Figure 4.4: The Data-Stall Frontier: In the ideal overlapped pipeline, the compute active-time fraction is min(1, available bandwidth / input demand), plotted against a logarithmic bandwidth ratio. Markers place three of the chapter's workloads on the curve: the data-stall notebook's pipeline, whose delivery time exceeds its compute window, and the image and text jobs of the text-versus-image bandwidth comparison, each served by the whole parallel file system; arrows from the knee measure each fed job's bandwidth headroom. The curve is an ideal service-rate bound, not an empirical utilization curve or a claim about a storage technology.

Below matching supply, bandwidth increases raise the ideal active fraction proportionally. Above it, more bandwidth does not improve this ideal compute stage, although headroom can help absorb contention and variable delivery. The markers show how far apart the chapter's workloads sit. Served by the whole 1 TB/s parallel file system, the text job of notebook 4.1 has about 6,000× more bandwidth than it consumes, the image job only 3.3×, and the pipeline of notebook 4.5 sits left of the knee at 80 percent active time. Real systems can fall below the curve because of preprocessing, latency tails, or synchronization. The graph identifies a workload service-rate boundary; it does not establish that every job needs every storage tier.

Because $\text{BW}_{\text{required}}$ grows linearly with N_{GPU}, adding compute nodes against a fixed shared file system moves any workload leftward along figure 4.4 until supply equals demand; past that saturation knee, additional GPUs add cost without raising training throughput, and equation 4.2 locates that limit from the storage tier's measured bandwidth. Even to the right of the knee, however, the frontier in figure 4.4 assumes that storage reads and CPU preparation overlap computation completely and that latency spikes never empty the input queue. Achieving that overlap in practice requires pipelining the loader stages and sizing a prefetch buffer deep enough to absorb tail delays.

4.4.1 Pipelining and prefetching

When a data loader waits until the GPU finishes batch k before reading and decoding batch $k+1$, the accelerator sits idle across the entire storage read and CPU preparation window on every step, as panel (a) of figure 4.5 shows. Pipelining removes that synchronous stall by staging and preparing batch $k+1$ on host resources while the GPU trains on batch k, turning the sum of input and compute time into their maximum (section A.2.1). When delivery keeps pace and enough prepared work sits in host memory to absorb latency variation, four batches finish in about the time sequential loading

needs for two (panel (b)); at fleet scale, the open question is how deep that buffer of prepared work must be.

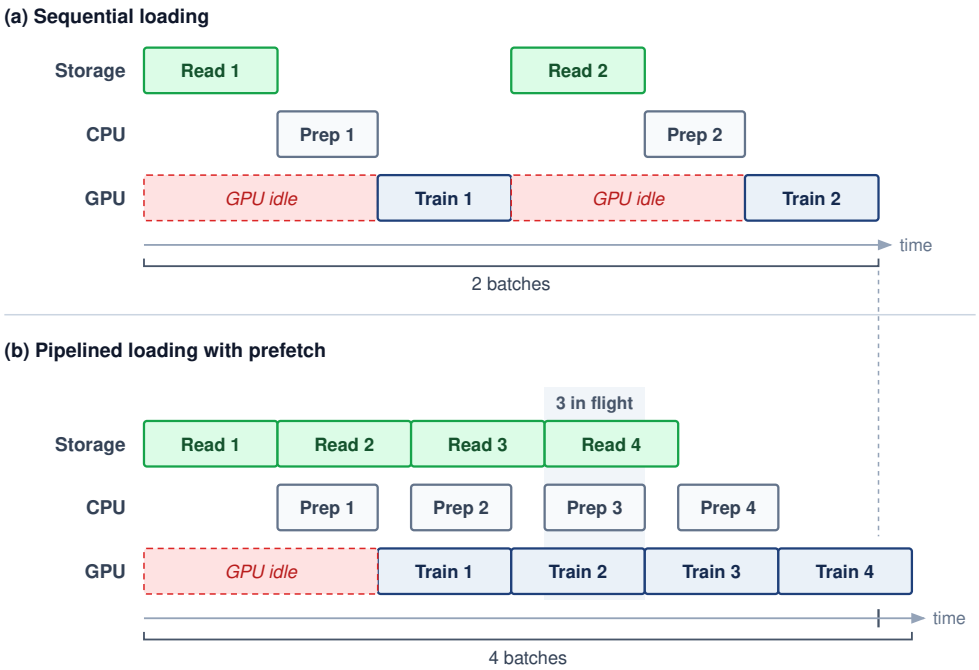

Figure 4.5: Pipelined Data Loading: Both panels share one time scale, with illustrative stage durations. Sequential loading (a) leaves the GPU idle while each batch is read and prepared. Pipelined loading (b) reads and prepares batch $k+1$ while the GPU trains on batch k, so four batches finish in about the time sequential loading needs for two (dashed line). The shaded column marks one instant with three batches in flight. The overlap hides storage latency only while delivery keeps pace with compute and the prefetch buffer is deep enough to absorb variance.

Pipelining converts a sequential bottleneck into overlapped execution (figure 4.5), but continuous saturation requires a prefetch buffer deep enough to absorb variance in I/O latency. Stage durations need not be deterministic if ready batches cover the delays. Storage systems experience I/O jitter: variation in read latency caused by contention and background work. NVMe drives suffer occasional latency spikes during internal garbage collection (when the controller reorganizes NAND flash blocks) or wear-leveling operations. Parallel file systems incur spikes when multiple training jobs contend for the same OSS nodes, when the MDS processes a burst of metadata operations, or when background scrubbing detects and repairs bit errors. Object storage likewise exposes delays from shared-service contention and maintenance.

A single slow I/O creates a "bubble" when it exhausts the pipeline's ready work. If the CPU is preparing batch $k+1$ with no additional batches ready, a delayed read can leave batch $k+1$ unavailable when the GPU finishes batch k. The GPU then idles until the batch is ready. A longer queue tolerates a longer interruption, and delivery faster than consumption can refill it afterward. Rather than preparing just one batch ahead, the CPU maintains a queue of N_{prefetch} preloaded batches.

A buffer sized to the measured p99 I/O delay needs enough ready batches to cover that delay without draining:

$$Q_{\text{prefetch,min}} = \left\lceil \frac{T_{\text{IO,p99}}}{T_{\text{compute}}} \right\rceil \tag{4.5}$$

The prefetch depth equation in equation 4.5 estimates how many ready batches cover the selected latency window. Outstanding reads contribute to replenishing that buffer but do not count as ready

batches until they complete. The equation assumes that the long-run data supply rate is at least as high as the accelerator's consumption rate; prefetching absorbs latency variance, not a sustained bandwidth deficit. When average throughput is sufficient but p99 latency exceeds compute time, deeper prefetching absorbs delays within the provisioned window; rarer or correlated delays can still cause stalls. When sustained delivery is lower than demand, identify the limiting stage before choosing more storage bandwidth, additional decoding capacity, a different representation, or local staging. The equation uses the p99 delay rather than the mean because a buffer sized for the mean drains whenever a slow read arrives, and the tail-latency arithmetic of section 4.3.6 shows that when a step depends on many servers, some read is slow on most steps. Even the p99 window leaves rarer events and correlated interruptions to test.

If object-storage reads take 500 ms at the 99th percentile and compute takes 200 ms, then $Q_{\text{prefetch,min}} = 3$ batches; provisioning five ready batches adds a margin. A data loader's configured work-ahead count (for example, a per-worker prefetch factor) is not a count of ready batches, because some of that work may still be loading or preprocessing. The quantity that protects the accelerator is completed-queue occupancy and its refill rate, measured against the delay window on whichever tier serves the reads.

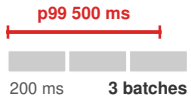

A 500 ms p99 read spans 2.5 steps of 200 ms, so three ready batches must cover it.

The image job of notebook 4.1 shows how prefetch buffer depth follows the serving tier. Each GPU processes its batch in roughly 200 ms, and once its 256 images are decoded into $224 \times 224 \times 3$-byte tensors, the collated batch occupies about 40 MB per GPU. Reading from local NVMe, the p99 I/O latency for that batch is approximately 50 ms. Under the p99-window model, the minimum ready-buffer depth is $\lceil 50/200 \rceil = 1$ batch, and this example provisions a depth of 2. Reading from a parallel file system, the p99 I/O latency rises to roughly 200 ms due to network jitter and contention, requiring a minimum depth of $\lceil 200/200 \rceil = 1$, with an illustrative provisioned depth of 3. Reading from object storage, the p99 latency can exceed 500 ms, requiring a depth of at least 3, with an illustrative provisioned depth of 5 or more.

These ready-batch depths translate directly into host DRAM consumption. At 40 MB per batch, a depth-5 prefetch buffer per GPU consumes 200 MB, and 8 GPUs per node consume 1.6 GB. At 40 MB per batch with a depth of 1, the same node needs only 320 MB. Measured delivery delays, batch size, and the chosen protection window determine this memory budget. Queue occupancy and refill rate must still be tested under correlated delays.

The representation decides whether that depth is cheap. Text barely registers, because a per-GPU batch of a few packed 4,096-token sequences occupies tens of kilobytes, so even a deep text queue is negligible. Decoded images are where depth costs memory. The 256-image batch above is roughly 39 MB before rounding, and a buffer of 8 such batches consumes about 308 MB per GPU, modest against a node's 512 GB of DRAM even after multiplying by the GPUs on the node. Higher-resolution images and video raise the per-batch size by an order of magnitude or more, at which point the prefetch buffer competes with embedding tables and other DRAM consumers.

Prefetching and the storage hierarchy together form a layered defense against stalls. The prefetch buffer in host DRAM absorbs I/O variance from NVMe reads. The local NVMe cache absorbs network variance from shared storage (whether a parallel file system or an object store). The parallel file system read-ahead cache absorbs variance from disk seeks across the OSS cluster. Because each layer adds latency tolerance at the cost of memory or storage capacity, the system designer must check each active path against the interruption window it is intended to absorb and verify that sustained delivery replenishes the consumed batches.

The image job shows what staging buys when steps are short. Its 500 ms object-store p99 against a 200 ms step calls for 3 ready batches, while a local NVMe hit at 50 ms needs only 1. The 175B running example sits at the other end of the scale. Against its modeled 4.64-second step, even the object-store p99 is about 11 percent of one step, so $\lceil 500/4{,}639 \rceil = 1$ ready batch covers the p99 on every tier, and staging its shards on local NVMe does not reduce its prefetch depth. The running example stages for other reasons, which section 4.8 weighs. The copy happens once before training starts, and local reads remove the rarer and correlated object-store delays that a buffer sized to the p99 does not cover.

Checkpoint 4.3: Data pipeline design

A training cluster the size of the book's reference cluster runs 1,024 GPUs with 128-image batches (116 KB per image after compression, 150 ms per iteration).

- ☐ **Bandwidth requirement**: with the pipeline bandwidth requirement of equation 4.2, which scales accelerator count and target utilization by bytes per iteration, calculate the required aggregate storage bandwidth at 90 percent target utilization.
- ☐ **Checkpoint contention**: determine whether 800 GB/s aggregate parallel-file-system bandwidth is sufficient, then recalculate when checkpoint writes consume 20 percent of that bandwidth.
- ☐ **Prefetch depth**: calculate the minimum prefetch buffer depth in batches needed to absorb 300 ms p99 object-storage latency.
- ☐ **Host memory**: calculate total prefetch buffer memory when each decoded batch occupies 10 MB in host DRAM at the depth calculated in (3).

4.4.2 Multi-worker data loading

A single CPU core cannot keep a high-throughput accelerator fed, even when the storage hardware is fast enough, because of the CPU work required between the storage read and the GPU transfer. For image training, decoding a 116 KB JPEG image into a $224\times224\times3$ raw tensor costs several milliseconds of CPU time. One measurement decoded full-resolution ImageNet images at about 530 per second across the four virtual CPUs of a cloud inference instance (Kang et al. 2020), and each augmentation, such as a random crop, flip, or color jitter, adds further work per image. The image job of notebook 4.1 feeds each GPU 256 images every 200 ms step, more than 1,000 images per second. At 1,000 images per second per GPU and 8 GPUs per node, a single core would need to decode and augment 8,000 images per second, tens of times what a single core can sustain.

The solution is multi-worker data loading, where W worker processes each read from storage, decode, augment, and enqueue batches independently. The effective I/O throughput scales approximately linearly with W until one of three bottlenecks is reached: the storage device's bandwidth saturates, the PCIe bus between host and device saturates, or the workers exhaust host CPU cycles.

The worker count W is tuned as on a single node, by profiling the data stall ratio, since too few workers starve the GPU and too many waste CPU on contention. At fleet scale, every core given to loading is a core the training runtime, communication threads, and operating system cannot use, so dividing host CPU between data loading and training orchestration is itself a capacity planning exercise. How those workers exploit concurrency then depends on which tier serves their reads. On local NVMe, workers issue concurrent reads to the same RAID array and rely on NVMe's deep command queues to expose device parallelism until shared drive or controller capacity saturates. On a parallel file system, they distribute reads across different OSS nodes to aggregate server bandwidth. On object storage, they issue concurrent HTTP GET requests to distinct shards.

Multi-worker loading must also preserve shuffle quality without giving up I/O efficiency. Under the shard-level and within-shard shuffling of section 4.2, the loader assigns contiguous groups of the shuffled shard list to each worker, and the adequacy of that mixing depends on how the shards were built. A corpus whose shards are internally homogeneous, such as one source or one class per shard, needs finer shards or a larger shuffle buffer, so the layout should be validated against the loss curve of a fully shuffled baseline on a subset before the run commits to it.

The shuffle buffer size creates a memory-randomness trade-off. A larger buffer provides better randomness (because samples are drawn from a larger pool) but consumes more host DRAM. For the running example with text data, a shuffle buffer of 10,000 sequences at 2,048 tokens each, with each token represented as a 4-byte integer, consumes roughly 82 MB. This is negligible relative to the 512 GB of host DRAM. For image data, where each decoded image occupies 151 KB, a shuffle buffer of 10,000 images consumes 1.5 GB, still manageable but a nontrivial fraction of the prefetch budget. The data loader designer must balance shuffle buffer size against prefetch depth, since both compete for the same host DRAM capacity.

The data loading pipeline, from storage read through decode, augmentation, and transfer to accelerator, is the complete fuel delivery system for the training engine, and when any stage becomes the bottleneck the accelerator starves. Local reads can reduce network and shared-storage wait, which makes shard placement across nodes the next decision when that path limits delivery.

4.4.3 Data locality and placement

Even when a node has enough CPU workers to decode batches at full speed, fetching every shard across a shared storage network exposes the training loop to fabric congestion and remote service queues. Placing shards on node-local NVMe avoids that shared-network contention, whereas keeping them in a remote parallel file system or object store preserves a single shared pool that any node can read without copying data ahead of time. Because local NVMe's advantage depends on both the node's drive stripe and the per-client bandwidth the remote tier can sustain, placement decisions must compare measured client rates rather than assume a fixed local-to-remote ratio.

The staging decision determines where the job pays for data movement, whether before training starts, during the first epoch, or on every epoch. The simplest strategy is static placement, or *prestaging*. At the beginning of a training job, the orchestrator assigns a fixed subset of data shards to each node, and the node copies its shards from shared storage (the object store in the running example, a parallel file system in many on-premises clusters) to local NVMe drives before the first iteration. Reads remain local while the assigned shards fit and the sampling policy uses that assignment, which gives the highest steady-state throughput at the cost of the longest startup delay. Staging a 10 TB dataset across 256 nodes at 500 MB/s per node takes roughly 80 seconds under ideal conditions, and contention on shared bandwidth can extend that stall to several minutes. A changed dataset or assignment requires restaging, so the placement policy must preserve the intended sample distribution as well as local throughput.

The alternative is dynamic placement, or *on-demand staging*, where nodes fetch shards from shared storage when the data loader first accesses them and retain them locally for reuse. This removes the upfront delay and supports datasets larger than aggregate local NVMe capacity and sampling policies that change the active subset over time, but the first epoch pays the full network-read cost. Cache misses consume shared-storage bandwidth and expose remote latency unless prefetching hides it; cache hits retain the benefit of local reads, so the pattern pays off when later epochs reread the cached shards.

The cache view of equation 4.1 puts a number on that trade. Take the multi-modal job of section 4.3.4, whose 35 TB corpus overflows a node's 30.7 TB of local NVMe. If the job samples uniformly and the cache holds that fraction of the corpus, about 88 percent of reads hit the 28 GB/s local stripe and the remaining 12 percent fall through to the node's 4 GB/s share of the parallel file system. The delivered rate of a serial loader is then a harmonic mean, about 16 GB/s, only 58 percent of the all-local rate. The service time here is the time to deliver one byte, so the loader averages the time per byte across hits and misses rather than averaging the two bandwidths directly. A miss rate near one in eight costs about two-fifths of the bandwidth because each missed byte takes 7× as long as a hit, and that mean is the optimistic case; misses that cluster in time drain the prefetch buffer faster than the average suggests.

Streaming keeps data off local storage entirely, reading records from the parallel file system or object store on every epoch. It has no startup delay and needs no local NVMe capacity, but every epoch absorbs the network-read cost again. On transfer cost alone, streaming suits single-epoch pretraining, which reads each byte once, and datasets that exceed the cluster's aggregate local capacity. The single-epoch running example stages anyway, for reasons of latency tails rather than transfer fees, as section 4.8 shows.

A more advanced approach is locality-aware scheduling. The scheduler keeps track of which data shards are cached on which nodes' local NVMe drives. When it starts a new job or replaces a failed node, it prefers a node that already has the required data cached, rather than treating all nodes as interchangeable. This uses the principle of data gravity,[11] co-locating compute with data to minimize transfer times. Scheduling avoids a staging copy only when an available node holds the required, validated shard versions. The next requirement is therefore to identify the cached contents precisely enough to reuse them safely.

[11] **Data gravity**: A metaphor treating large datasets as massive objects that attract compute. The gravitational analogy is apt: moving a 100 TB dataset out of one cloud provider to another costs roughly $9,000 in egress fees and takes hours, while launching compute next to the data is near-instantaneous. In ML fleet design, data gravity means the storage location of the training dataset often dictates the physical placement of the entire training cluster.

For the 6 TB serialized dataset distributed across 256 nodes, each node is responsible for approximately 23.4 GB of data. If a node fails, its replacement must restage that shard subset from shared storage. Drawn at the uncontended 4 GB/s per-node share of the chapter's parallel file system, the copy takes about 5.9 seconds. At an effective 400 MB/s per-node rate, which recovery contention can impose and which is also the order of the per-node share of the 50–100 GB/s object-store streaming path described in section 4.3.7, the same copy stretches to about 58.6 seconds. If the orchestrator can instead schedule the replacement workload on an idle node that already has the data cached from a previous run, staging time is zero and training resumes as soon as the replacement joins. The saving is modest per event but recurs. At a per-GPU mean time between failures of 50,000 hours, GPU faults alone give the 256-node cluster about 1 node failure per day, and each contended restage idles the whole synchronous job for the copy time above. Larger fleets fail proportionally more often, although each node's share of the corpus shrinks as the node count grows.

The pipeline equation, the prefetch queue, the loader workers, and shard placement together determine which bytes reach host memory and when. On the conventional buffered path, however, those bytes are also copied by the CPU through host DRAM before the accelerator's DMA can start, and at high request rates that copy competes with decoding for the same cores. Orchestration does not remove that copy; a more direct data path between storage and accelerator can.

4.5 GPUDirect Storage and the CPU Data-Copy Bypass

Even when the storage tier has enough raw bandwidth, the traditional data path for loading training data routes every shard read from storage → host DRAM → GPU HBM in three hops, one of them a CPU copy inside host DRAM. Data is first read from NVMe into a kernel buffer in host DRAM by a DMA transfer, then copied by the CPU into the data loader's user-space buffer, and finally transferred over PCIe to GPU memory by the GPU's DMA engine. Each hop adds latency and consumes host resources because the CPU must orchestrate every transfer, manage buffer allocation, and handle interrupts from both the storage device and the GPU. For a single transfer, this overhead is negligible, but when eight GPUs on a node each demand thousands of small transfers per second, the aggregate CPU load for payload movement alone can saturate multiple cores and starve the data augmentation stage. GPUDirect Storage (GDS),[12] which removes that host-memory bounce copy for local or RDMA-attached NVMe while the CPU still submits and coordinates I/O through the software stack, pays off whenever that host copy limits throughput.

[12] **GDS (GPUDirect Storage)**: Part of NVIDIA's GPUDirect family, which also includes GPUDirect RDMA (network adapters writing directly to GPU memory) and GPUDirect Peer-to-Peer (inter-GPU memory access); GDS extends this data-copy-bypass approach to NVMe storage. With RDMA and GDS, payload data can traverse the path remote storage → network → GPU memory without a CPU-mediated bounce copy, while the CPU still issues and coordinates the operations.

Definition 4.4: GPUDirect Storage

GPUDirect storage (GDS) is a technology that enables a direct DMA path between NVMe storage devices and GPU memory, bypassing the system-DRAM bounce buffer and CPU payload copies while retaining CPU submission and control.

1. **Significance**: It eliminates the "bounce buffer" through system memory. In the small-transfer model detailed in notebook 4.6, the assumed active CPU service per request falls by 75 percent. That model estimates CPU capacity savings, not an end-to-end latency guarantee. Whether the storage link saturates depends on the remaining bottlenecks.
2. **Distinction**: Unlike traditional I/O, where payload data is staged through host buffers and copied between them, GDS provides direct memory access between the storage controller and the accelerator. CPU-issued I/O calls and kernel storage drivers remain in the control path.
3. **Common pitfall**: A frequent misconception is that GDS makes all storage "faster." The direct path depends on supported storage and interconnect configurations; it does not eliminate media latency, network latency, or queueing.

The software layers in the traditional I/O path contribute control overhead as well as payload-copy work. Direct I/O and pinned buffers can avoid some copies even without GDS. The correct baseline is therefore the actual data path rather than a fixed number of copies for every read. The direct path changes the payload route without eliminating the software control path. The application registers a GPU memory region for direct DMA (through the cuFile API in NVIDIA's implementation), and

later reads move payload from NVMe into GPU memory without an intermediate system-memory buffer. The file system and kernel storage drivers still handle submission and completion.

The benefit from GDS matters when host payload movement is a limiting stage after other pipeline work is accounted for. In figure 4.6, the conventional path takes two DMA hops through intermediate host-DRAM bounce buffers, while GDS takes a single direct DMA transfer. The direct path eliminates the CPU payload-copy stage, but the dashed control link shows that the host CPU still orchestrates transfer submission and completion.

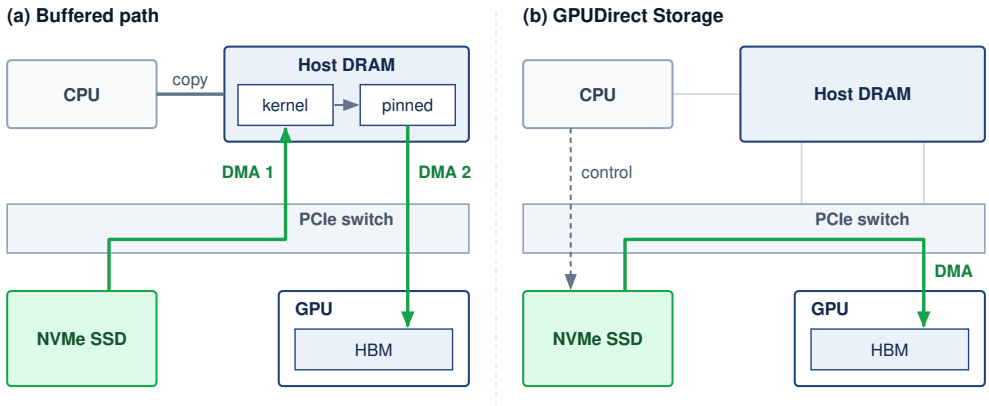

Figure 4.6: Traditional vs. GPUDirect Storage Path: Both panels draw the same node: CPU and host DRAM above a PCIe switch, NVMe and GPU below it. Green routes are DMA transfers of payload bytes. In the buffered path (a), payload climbs into host DRAM and comes back down, and the CPU copies every payload from the kernel buffer to a pinned buffer before the accelerator's DMA can start. In the GDS path (b), NVMe transfers directly into GPU memory, and the dashed CPU link shows the submission and completion work that remains.

GDS benefits depend on the bottleneck, request sizes, alignment, concurrency, and storage topology. Removing host bounce buffers can improve throughput when those copies limit the path; it does not imply that the smallest transfers benefit most. Small I/O may remain submission- or latency-limited and require concurrency to approach the available bandwidth. Benchmarks of the direct path therefore measure I/O size and per-device behavior as well as throughput. The benefit appears only in an end-to-end comparison of the pipeline with and without the direct path, and the clearest case is the one where host CPU time per request is the limiting stage.

Napkin Math 4.6: The CPU bypass dividend

Problem: A training node with 8 GPUs loads 116 KB images at 8,000 images/s per GPU (64,000 images/s total). That rate is deliberately aggressive, several times the per-GPU image rate of notebook 4.1, to isolate the case in which per-request CPU work is the limiting stage. Assume active CPU service of 120 µs per request on the buffered path and 30 µs on the direct path. These illustrative service times exclude time blocked on I/O; compare the implied CPU load rather than treating elapsed transfer latency as CPU time.

Traditional path: In the assumed buffered implementation, each image requires a DMA from NVMe to DRAM, a host copy from kernel to user space, and a PCIe transfer to GPU. At 64,000 images/s with 120 µs of CPU time per image, the CPU spends 7.68 seconds of CPU time per wall-clock second, consuming roughly 8 cores worth of processing just for data movement.

GDS path: Each image is transferred directly via DMA from NVMe to GPU. At 30 µs of CPU time per image (for initiating the DMA), the CPU spends 1.92 seconds of CPU time per wall-clock second, freeing roughly 6 cores for data augmentation.

Systems insight: Under the assumed CPU-service costs, bypassing payload copies releases processing capacity for preprocessing. The calculation does not predict end-to-end transfer latency, nor prove that preprocessing or storage throughput can sustain this request rate.

GDS has practical limitations. Not all NVMe controllers support peer-to-peer PCIe transfers, and the data must already be in a format that the GPU can consume directly (raw pixel values, token IDs, or formats handled by GPU decoders, rather than arbitrary compressed JPEG or gzip streams). In practice, many pipelines use a hybrid approach: CPU-side decoding for compressed formats followed by a normal pinned-memory GPU transfer, while GDS is reserved for direct reads of GPU-consumable data from supported local or remote storage.

The architectural evolution from CPU-mediated copies to direct storage access reflects a broader trend of removing CPU payload copies from the critical data path wherever possible. RDMA allows network payloads to move without CPU copying, while GDS provides the same bypass for storage payloads. Together, RDMA and GDS let data flow from remote storage (such as NVMe over Fabrics, or NVMe-oF), through the network fabric, and into GPU memory without a host-memory bounce copy. The CPU still submits and coordinates the operation, while more of its capacity remains available for data augmentation, pipeline orchestration, and error handling.

↩ RDMA and kernel-bypass network paths are established in section 3.4.1.

The GDS design principle also extends to checkpoint writes. Whereas a host-staged path transfers checkpoint bytes from GPU HBM through host DRAM to NVMe, a supported direct path can bypass that host staging for a 2,450 GB checkpoint. Removing the copy reduces T_{write} only when that movement contributes to the exposed pause; snapshot consistency, serialization, storage service, and durable publication remain separate requirements.

4.5.1 The complete data path

Combining GDS with the multi-tier hierarchy, the complete data path connects the tiers for the running example. At the beginning of the training job, the 6 TB serialized shard set, derived from the 3 TB compressed source corpus, is staged from object storage into the cluster's local NVMe caches, with each node receiving its assigned shard subset during setup. The data loader workers read tokenized shards from NVMe, assemble packed sequences in pinned host DRAM, and transfer those batches to the accelerator. With GDS enabled, token IDs that are already in a GPU-consumable layout can be loaded directly into GPU memory without CPU payload copying; compressed formats still require CPU or GPU decompression before use.

During steady-state training on the host-staged path, CPU data loader workers read from local NVMe and fill a prefetch queue in host DRAM. The GPU receives batches from this queue via PCIe DMA; the supported GDS alternative bypasses that host payload staging. The compute phase (forward pass, backward pass, optimizer step) consumes the batch and updates model state in HBM. Every 10 minutes, the training framework captures a consistent checkpoint, writes it to local NVMe through the chosen path, and copies it to the parallel file system in the background. The checkpoint protocol must still establish when all required shards are safe to publish for recovery.

At the end of the training job, the final model checkpoint is promoted from the parallel file system to object storage for long-term retention. The local NVMe copies are deleted during job cleanup. If the model is deployed for inference, it is loaded from object storage into the serving cluster's HBM, completing the lifecycle of data movement through the hierarchy.

Training samples and checkpoints are each copied several times in this lifecycle. Each training sample traverses one chain of copies: object storage $\rightarrow$ NVMe (staging), NVMe $\rightarrow$ host DRAM (read), host DRAM $\rightarrow$ GPU HBM (transfer). Each checkpoint traverses another: GPU HBM $\rightarrow$ NVMe (local save), NVMe $\rightarrow$ parallel file system (async copy), parallel file system $\rightarrow$ object storage (long-term retention). Every copy consumes bandwidth and contributes latency. The engineering goal is to minimize copies on the critical path (the training loop) and tolerate additional copies on noncritical paths (staging and archival). Both lifecycles cross the same tiers in opposite directions, with synchronous critical-path copies drawn solid and asynchronous background transfers drawn dashed in figure 4.7. For training inputs moving rightward, the active loop waits only on local promotion into HBM (or the optional GDS bypass); for checkpoints flowing leftward, the exposed pause is confined to the local NVMe save while parallel-file-system replication and archival retention proceed off the critical path.

Every copy that leaves the node also needs a network, and the choice of network decides whom a storage burst can disturb. Many training nodes carry a small number of storage-facing network adapters on a separate front-end network, in addition to one compute-fabric adapter per accelerator. Routing staging reads and checkpoint copies over that front-end network isolates them from gradient

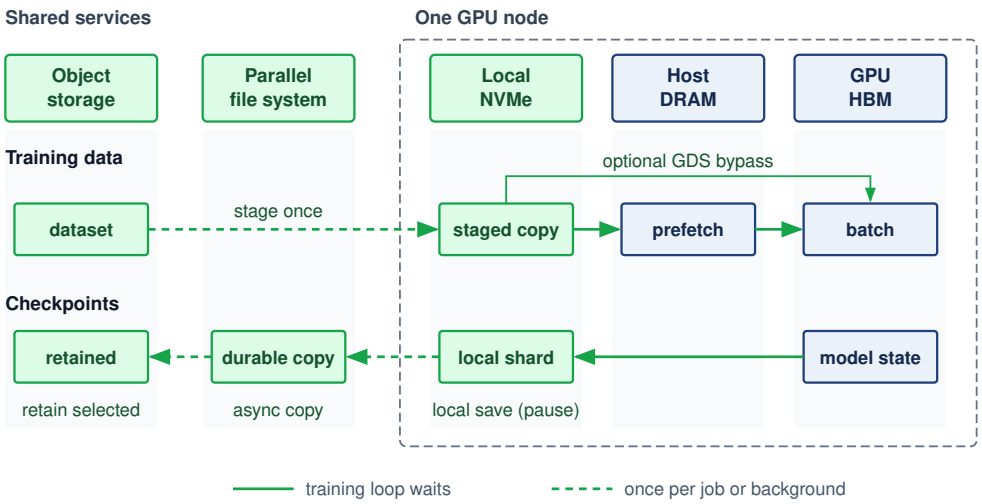

Figure 4.7: The Complete Data Path: Columns are storage tiers, ordered from the durable object store (left) to accelerator memory (right). Training data moves right and bypasses the parallel file system in this running example; checkpoints return left through it. Solid arrows mark copies the training loop waits on; dashed arrows mark copies that run once per job or in the background. The dashed outline is one GPU node: every copy the training loop waits on stays inside it.

collectives, but it caps each node's storage bandwidth at a few front-end ports. Routing storage over the compute fabric offers far more bandwidth, and remote GDS reads use exactly that RDMA path, but then a checkpoint burst shares switch buffers and spine links with the job's collectives and must be separated by traffic classes.

For the running example, each node's fair share of a 1 TB/s file system shared by 256 nodes is roughly 4 GB/s, and copying its roughly 10 GB checkpoint shard within the 10-minute interval needs far less, so a front-end network can carry it. Image workloads that stream hundreds of GB/s across the cluster push designs toward the compute fabric. Through the C^3 lens, storage traffic becomes communication (C_2) load the moment it leaves the node, and the network chosen to carry it decides which collectives a storage burst can delay.

Over the 30-day window, checkpointing writes roughly 10.6 PB of state against a single read of the 6 TB serialized dataset, so checkpoints, not training inputs, are the storage workload to design next. That design must take a multi-terabyte save off the training critical path without ever exposing a checkpoint that cannot be restored.

4.6 Checkpoint Storage

After the 2,048 accelerators of the running cluster have been training for six hours, a power supply failure turns checkpoint storage into the difference between resuming within minutes and repeating hours or days of lost computation. Whether the job recovers cleanly depends on three storage questions: how quickly the most recent checkpoint was saved, where its shards reside, and whether they can be read back intact. Because checkpoints are also the heaviest write workload in the hierarchy, the storage architecture must minimize T_{write}, the time the training pipeline pauses to save a snapshot, while stating when that snapshot becomes durable and visible for restore. The save interval is an input to this storage analysis; selecting it from failure rates and recovery costs is a separate reliability decision derived in section 7.8.1.

Under the resumable mixed-precision Adam contract of section 4.1, the 175B parameter model's 2,450 GB checkpoint comprises 350 GB of FP16 model weights, 700 GB of FP32 master weights, and 1,400 GB for the two FP32 Adam moments, while learning-rate-scheduler state, random-number-generator state, and the data-loader cursor add negligible bytes and gradients are recomputed after restore. When a save triggers, every participating worker process, called a rank, writes its assigned shard at once, producing the checkpoint storm introduced in section 4.2.

Keeping that burst tractable requires the write-once rule, under which a coordinated save writes each unique byte once, whatever the parallel layout. In the running layout of table 4.3, each node holds the 1/8 of a replica's state that belongs to its pipeline stage, and all thirty-two replicas hold identical copies of it. A coordinated save therefore assigns each replica a disjoint 1/32 of every stage slice, so each of the 256 nodes writes about 1/256 of the checkpoint. The third ZeRO stage (ZeRO-3), which partitions the weights as well, makes that even split the natural one, because every rank already owns a distinct slice.

Breaking the write-once rule is a real risk even in the running layout, where a save in which every replica writes the full state it holds would multiply the checkpoint bytes by the replica count. Its cost is clearest in a deliberately extreme thought experiment, a smaller model whose complete resumable state is written once per GPU of the book's reference cluster.

Writing one copy per replica takes 1,024 times longer than writing each unique byte once.

Definition 4.5: Checkpoint storm

Checkpoint storm is a burst of synchronized network and storage traffic that occurs when all nodes in an ML training fleet save model state simultaneously.

1. **Significance**: When each replica writes its own full copy, the naive storm takes $T_{\text{write}} = N_{\text{DP}} \times \text{per-replica state}/\text{BW}_{\text{ingest}}$, where N_{DP} is the number of replicas writing at once and $\text{BW}_{\text{ingest}}$ is the aggregate rate at which the target storage tier absorbs writes. A coordinated save that writes each unique byte once takes unique state/$\text{BW}_{\text{ingest}}$. Under the chapter's checkpoint contract, a 70B-parameter model, smaller than the 175B running example, stores 140 GB of FP16 weights, 280 GB of FP32 master weights, and 560 GB of Adam moments. With vanilla data parallelism across 1,024 replicas (one per GPU of the book's reference cluster, rather than the chapter's 256-node fleet), naively writing that complete state from every replica generates 1003.5 TB of simultaneous traffic. No 80 GB device can hold a 0.98 TB replica, so the case is a bound on write volume rather than a deployable layout. At 100 GB/s of aggregate storage ingest bandwidth, the write takes roughly 10,035.2 seconds, or over 167.3 minutes of training stall. A coordinated checkpoint protocol writes one logical copy. Replicated data parallelism can select one writer; ZeRO-3 instead spreads that copy across the ranks, each writing the slice it owns. In the sharded example the total is roughly 0.98 TB, bringing T_{write} down to roughly 9.8 seconds. Gradients do not appear in either total because recovery recomputes them.
2. **Distinction**: Unlike general I/O contention (which is stochastic and unpredictable), a checkpoint storm is synchronous and periodic: every node writes at the same moment because the training orchestrator triggers checkpoint after a fixed step count. Its predictability makes it both more damaging (all nodes compete simultaneously) and more tractable (it can be prevented by design through staggered scheduling or asynchronous serialization).
3. **Common pitfall**: A frequent misconception is that checkpointing every 100 steps is "low overhead" whatever the save protocol. Under the naive protocol at 70B scale, if each training step takes 10 seconds and a checkpoint storm takes 10,035.2 seconds, checkpointing every 100 steps means spending most of the run on checkpoint I/O rather than useful training.

The severity of a checkpoint storm depends on two factors: the per-node shard size, determined by the model size and the sharding strategy, and the throughput of the first write tier, typically local NVMe. A concrete 175B example shows why the local-write phase, rather than the durable copy to the parallel file system, governs the exposed training pause.

Napkin Math 4.7: Tiered checkpoint staging

Problem: Assume sustained bandwidth, local drives that write as fast as their rated sequential read speed, no snapshot-copy, serialization, or coordination overhead, and no interference from background copies. The resulting times are bandwidth lower bounds. A 256-node cluster saves a 175B-parameter checkpoint every 10 minutes. Each checkpoint totals 2,450 GB. Writing each unique byte once and spreading the bytes evenly across nodes, each node saves roughly 9.6 GB.

1. **Per-node write to local NVMe** (4 drives at 7 GB/s each = 28 GB/s): 9.6 GB ÷ 28 GB/s ≈ 0.34 seconds.
2. **Async copy to PFS**: 256 × 9.6 GB ≈ 2.4 TB total. If the PFS provides 1 TB/s aggregate, the storm completes in roughly 2.4 seconds.
3. **Per-node PFS bandwidth**: If all 256 write simultaneously, each gets 1000 GB/s ÷ 256 = 3.9 GB/s.
4. **Training pause**: At least 0.34 seconds for the local NVMe write under this model. The PFS copy overlaps with the next training iteration.
5. **Overhead lower bound**: 0.34 seconds pause every 600 seconds ≈ 0.06 percent training time lost to checkpointing.

Systems insight: Under these 1 TB/s aggregate-PFS assumptions, tiered staging reduces checkpoint overhead from about a 2.4 seconds PFS-direct write to a sub-second local write. The training pause can end at the local snapshot boundary before durable publication, but actual pause includes snapshot creation, serialization, and coordination in addition to this byte-transfer lower bound.

A PFS-direct save of a few seconds every 10 minutes costs well under 1 percent of the run, which can make tiered staging look unnecessary for the running example. The checkpoint-dominance argument of section 4.5.1 concerns bytes moved, not this pause, and the notebook's few-second save holds only while its assumptions do. The job has the whole 1 TB/s file system to itself; shared among the ten concurrent jobs of section 4.3.6, it receives a tenth of the bandwidth, the same save takes ten times as long, and its burst slows the other jobs' reads. Realized ingest also falls below the rated aggregate once serialization, metadata operations, and stragglers enter. The save interval shrinks as the fleet grows, because a larger fleet fails more often and the reliability analysis of section 7.8.1 then calls for more frequent saves, and a larger model scales every save with it.

Tiered staging insulates the training loop from that shared-tier contention by splitting the save into two phases. In the first phase, as calculated in notebook 4.7, each of the 256 nodes writes its 9.6 GB share of the 2,450 GB checkpoint across 4 local NVMe drives at 28 GB/s, establishing a byte-transfer lower bound on T_{write} of roughly 0.34 seconds before snapshot-copy, serialization, and coordination overhead. Once every required local shard has completed and the captured state is no longer being modified, the training pipeline can resume.

In the second phase, a background process asynchronously copies the local checkpoint to the parallel file system for durability. The copy can overlap training, but it still competes for memory, network, and storage service and must drain before repeated saves create an unbounded backlog. The two phases therefore land differently in the Fleet Law (equation 1.3). The local write is a pause that every rank takes at the same step boundary, and training resumes only when the slowest writer has finished its shard. That is the shape of a barrier rather than of a transfer between workers, a coordination (C_3) cost, so the pause adds to the synchronization term $T_{\text{sync}}(N)$ of each saving step, the term that already carries barrier and straggler waiting. The background copy belongs to no step's critical path. While it stays hidden and drains before the next save, it adds nothing to the step time; it reappears only as interference, in $T_{\text{comm}}(N)$ when it shares links with collectives or in exposed input wait when it slows shared-storage reads. Rate-limiting the copy, or staggering its start across nodes, spreads the burst over time and keeps that interference low, at the price of a later durability boundary.

If a node fails before its copy completes, the new generation may lose a required shard. Replication to peer nodes' NVMe drives can protect against a single-node loss, provided replicas occupy independent failure domains. Waiting for that replication extends the pause beyond the local-only calculation. A design that resumes earlier must retain the previous committed generation until the new one reaches its declared durability boundary.

Durability alone does not make a distributed checkpoint restorable. Each rank writes its checksummed shard under an immutable generation identifier. Only after every shard has persisted does a coordinator publish a manifest naming the complete shard set. That commit uses an atomic primitive of the target tier, an operation that takes effect completely or not at all. A file system offers a create or rename. An object store offers a conditional create, which succeeds only if the name is still unused, or a versioned pointer, a small record that names the current generation. A crash can then leave garbage to collect but never an incomplete generation that appears valid.

↪ The commit rule behind this ordering, a two-phase commit over the shard set, and the consistency of the cut that a manifest names are developed in section 7.8.5.

The exposed pause and the durability boundary therefore fall at different moments, as the four concurrent lanes of figure 4.8 align them. Training pauses only during the local NVMe shard write (t_0 to t_1), resuming on live state while the parallel file system absorbs the asynchronous copy (t_1 to t_2) and the coordinator validates checksums before publishing the manifest last (t_3).

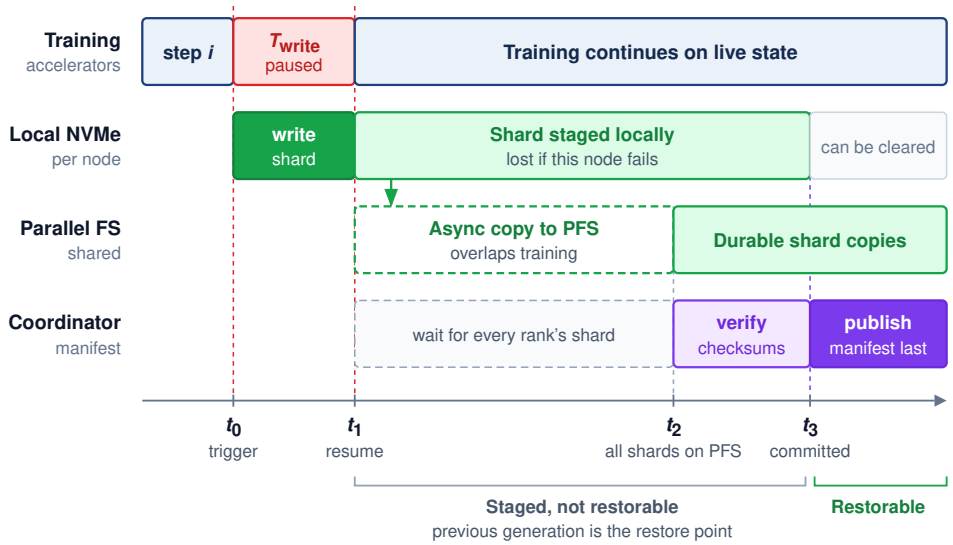

Figure 4.8: Two-Phase Checkpoint Staging: Read the lanes left to right. Training pauses only for the local NVMe write (t_0 to t_1); the copy to the parallel file system overlaps later training steps, and the coordinator publishes the manifest last (t_3). Between t_1 and t_3 the new generation exists only as staged shards, so the previous committed generation remains the restore point.

4.6.1 Synchronous and asynchronous checkpoint writes

In the two-phase timeline of figure 4.8, the first phase is still synchronous with respect to local flash. Across 256 nodes, every node must write and acknowledge its shard of the same completed training step before the completion barrier releases the next iteration, with the commit rules that make that cut consistent developed in section 7.8.5. The slowest local writer therefore sets the exposed pause from t_0 to t_1.

Asynchronous checkpointing shrinks even that local-flash pause by staging the snapshot in host DRAM first. Nodes coordinate the training step to capture, copy each accelerator-resident shard into a dedicated pinned host-memory buffer, and resume training immediately while a background thread writes the buffer to local NVMe. The snapshot buffer must remain immutable until the background write finishes; separate live-state buffers allow training to advance without changing the captured generation. If a node fails during the asynchronous write, its local shard may be incomplete. The job must wait for the required writes or peer replication before declaring this generation durable, even if training has already resumed.

The choice between synchronous and asynchronous checkpointing depends on the pause budget and save frequency. A 1-second pause costs the running example about 22 percent of one modeled 4.64-second step, or 0.17 percent of the 10-minute save interval. At that cadence a synchronous save is cheap; the overhead grows with pause length times save frequency, so the case for asynchronous checkpointing strengthens as saves come more often or pauses grow longer. Asynchronous checkpointing reduces exposed write time when snapshot buffers and background bandwidth are available, but its buffer capacity, contention, and durability delay also enter the decision.

Checkpoint formats implement the write-once rule of the checkpoint storm as rank-sharded saves, as in DeepSpeed's zero-redundancy checkpointing and PyTorch's Distributed Checkpoint (PyTorch Contributors 2026a). The sharding is what keeps the pause short. Had one node saved the full 2,450 GB itself, its 28 GB/s local stripe would need more than a minute rather than the 0.34 seconds of an even shard. Making each individual save fast, however, creates a cumulative capacity problem once thousands of snapshots accumulate over a multi-week run.

4.6.2 Checkpoint retention and restore

Checkpointing a 175B-parameter model every 10 minutes over a 30-day training window generates roughly 4,320 snapshots of 2,450 GB each, totaling approximately 10.6 PB, a volume that would crowd a shared parallel file system if retained in full. The retention policy this run adopts keeps the three most recent checkpoints on the parallel file system for fast recovery, copies every 100th checkpoint and the final checkpoint to object storage for long-term auditability, and deletes the rest. This policy reduces the parallel file system checkpoint footprint from 10.6 PB to roughly 7.35 TB (three live full checkpoints) while preserving 43 periodic historical snapshots in object storage, plus the final checkpoint when it does not coincide with that schedule, for posttraining analysis.

Incremental checkpointing offers a further optimization for reducing T_{write}. Rather than saving the entire training state every time, an incremental checkpoint records changed state relative to a declared base or preceding generation. The delta must include changed optimizer and continuation state as well as parameters. For models where a large fraction of parameters are frozen (such as fine-tuning scenarios where only the last few layers are updated), incremental checkpoints can be orders of magnitude smaller than full checkpoints. Restoring a generation requires its base and every intervening delta, so smaller saves add recovery work and dependencies. Periodic full checkpoints bound that chain, and the retention policy must preserve each dependency until no retained generation requires it. A shorter pause matters only if the checkpoint can be read back; the costlier failure is a checkpoint that was written but never verified as restorable.

War Story 4.1: The backup that had not been restored (2017)

Context: On January 31, 2017, a GitLab.com site reliability engineer working a database replication incident accidentally wiped the primary PostgreSQL data directory (GitLab 2017).

Mechanism: All five layered backup paths failed simultaneously: S3 dumps were silently empty due to a version mismatch, Azure disk snapshots were disabled, replication was broken, LVM snapshots were six hours stale, and WAL archiving was unimplemented.

Impact: Six hours of production database writes, user comments, and repository metadata were permanently lost, requiring an 18-hour emergency restoration from a staging snapshot.

Fix: GitLab instituted automated daily backup restoration testing, unified PostgreSQL versioning across environments, and implemented automated WAL archiving to S3.

Systems lesson: Storage reliability is a restore property, not a backup property. A checkpoint, replica, or dump only counts if the team regularly proves that it can be located, validated, and restored under incident conditions. ML systems face the same vulnerability. A model checkpoint written to object storage or a snapshot of a feature store (the system that stores model input features) is useless if silent version mismatches or missing restore tests prevent resuming a multi-day training run after node preemption.

For the running job, a restore test reads a retained generation's shards back, checks them against its manifest, and resumes training from them.

↪ The restore drill, which tests this recovery path before an incident relies on it, is developed in section 7.6.2.

Checkpoint 4.4: Checkpoint storage design

A training cluster of 256 nodes saves a 175B-parameter model checkpoint every 10 minutes. Each checkpoint is 2,450 GB total, distributed across all nodes.

- ☐ **Per-node state**: calculate the per-node checkpoint size.
- ☐ **Local write**: calculate T_{write} for the local write when each node writes to local NVMe at 28 GB/s.
- ☐ **Async copy**: calculate the per-node write bandwidth and async-copy time when the parallel file system provides 1 TB/s aggregate and all 256 nodes write simultaneously.
- ☐ **Tiered staging**: explain why tiered staging (local NVMe first, then async PFS copy) is faster than writing directly to PFS.

With the checkpoint accounting now derived in full, the running example's storage demands can be consolidated against the same hierarchy. Placing the training dataset, live model state, checkpoints, and archive lineage side by side reveals which direction of traffic dominates the training window.

Systems Perspective 4.2: The 175B model's storage footprint

The running example, a 175B-parameter model trained on 256 nodes within a 30-day window, places data on every tier of the hierarchy (table 4.5).

Table 4.5: 175B-Parameter Training Storage Footprint: Volume and primary storage tier for each data category in a 30-day, 256-node training run.

Category	Volume	Primary tier
Training dataset (compressed source)	3 TB	Object Storage
Training dataset (tokenized, per epoch)	6 TB	Object Storage → NVMe cache → Host DRAM
Model weights (FP16)	350 GB	GPU HBM (distributed)
Master weights (FP32)	700 GB	GPU HBM (distributed)
Adam moments (two FP32 states)	1,400 GB	GPU HBM (distributed)
Single resumable checkpoint	2,450 GB	NVMe → PFS → Object Storage
All checkpoints (30 days, 10-minute interval)	10.6 PB	NVMe (transient) → PFS (recent) → Object (archive)
Archive (retained checkpoints + dataset)	~116.8 TB	Glacier

The total data moved through the hierarchy is approximately 10.6 PB of checkpoint data plus 6 TB per 1 epoch of training data reads. For a single-epoch language model training run, checkpoint I/O dominates data loading I/O by a factor of over 1,764.

Checkpoints therefore set the run's write budget, while the far smaller training corpus carries a different obligation. Because training quality depends on the exact samples delivered at every step, each byte of that corpus must remain traceable to the source version and transformation code that produced it.

4.7 Versioned Data and Commit Boundaries

If a multi-week training run suffers a sudden loss spike or a deployed model regresses, an engineering team that did not record which shard versions, preprocessing code, and shuffle seeds fed each step cannot tell a model bug from a silent change in the input data. A checkpoint manifest provides one

commit boundary per save. The data pipeline that built the training corpus has many, one at every stage that turns raw inputs into processed shards. The storage system must record both the exact bytes read and the preprocessing code, shuffle seed, schema, and integrity checks that produced them before the run began.

> Lighthouse 4.1: Archetype B (DLRM at Scale): Feature store latency
>
> While Archetype A (GPT-4/Llama-3) trains on static, versioned shards, Archetype B (DLRM at Scale) trains on live feature streams that change every second as users click and interact. Replaying an exact batch then becomes a point-in-time problem, which a feature store answers by keeping a retained, versioned feature log.

Even when the underlying shard files are immutable, a distributed loader produces a stream of consumed samples whose versions and order can shift across epochs during a 30-day run. Dataset versioning provides a "Git for data" record by snapshotting the cryptographic hashes[13] of every data shard, the exact preprocessing code used for tokenization, and the random seeds governing the global shuffle. With that snapshot in place, engineers investigating a mid-run divergence can replay the exact batch that triggered the loss spike and test whether the culprit was a corrupted sample or numerical instability in the optimizer.

At petabyte scale, the data pipeline has to make every transformation reproducible. Raw documents are cleaned, tokenized, packed into training examples, and sharded across storage tiers. Each stage writes an immutable artifact, and each artifact records the input version, code version, schema, and statistical checks it satisfied. Once those dependencies are explicit, the pipeline forms a directed acyclic graph: later artifacts depend on earlier artifacts, and no stage is allowed to rewrite its own history. When a training job starts, the cluster manager verifies that each worker's local NVMe shards match the data snapshot declared for the run. Treating the data pipeline as a graph of versioned artifacts rather than a simple file transfer makes silent data drift detectable and attributable to a specific stage, code version, or input snapshot. Immutability also pays a storage dividend. Because an artifact never changes once written, a manifest can name it by content identifier, a hash of its bytes, and many experiments can reference the same unchanged shards without copying entire datasets.

↪ Point-in-time correctness, the leakage it prevents, and the storage cost of retained feature history are developed in section 12.4.3.

13 **Hash cost (SHA-256):** Verifying the integrity of a 100 TB dataset using SHA-256 consumes approximately 55.6 CPU-core hours at a per-core hashing rate of 500 MB/s. The work parallelizes across shards, but it still competes with decoding for host cores and must read every byte, so the validation step is best pipelined with ingestion rather than run as a separate pass before the fleet starts.

4.7.1 Replay and commit boundaries in data pipelines

When a tokenization worker crashes halfway through building a shard of the running example's corpus, its replacement must decide which raw documents to process again without silently dropping records or writing duplicates. Versioned artifacts identify *what* a pipeline produced, while commit and replay rules determine whether mid-stage failures can alter that result. The storage system's part in that decision is the commit boundary, the point at which a stage's output becomes durable and visible downstream, the same contract that section 4.6 applied to checkpoint generations. That shared contract is why a storage chapter owns these semantics, because whether a replay loses or duplicates data is decided by what the store has committed. The restarted worker compares its input position with the last committed output and replays everything after it, so the work at risk is whatever accumulated past the last commit.

Two service-rate effects govern how much uncommitted work accumulates. Records are assigned to workers by a partition key, such as a document's source or a user identifier, and partition skew arises when one key owns disproportionate work and holds its stage open while the rest of the fleet idles. Backpressure propagates a slow downstream stage's limit upstream by slowing admission or bounding in-flight data, a control response to overload rather than another name for slow processing.

Failures near the commit boundary raise a correctness question, and three guarantees answer it differently. At-most-once processing never replays uncertain work, so a document that was read but whose output was never confirmed is lost. At-least-once processing replays everything after the last confirmed commit, so a document whose output was written but not yet confirmed is written twice. An effectively exactly-once outcome combines replayable inputs with idempotent writes, writes that leave the same result however many times they repeat, or with transactional output keyed by a stable event or transaction identifier (Akidau et al. 2015). Reliably delivering each record from

one worker to the next cannot provide this guarantee on its own, because the duplicate arises from the worker's own replay. The streaming feature pipelines of Archetype B (DLRM at Scale) face this question on every event.

> Example 4.2: The acknowledgment gap
>
> **Failure trace**: A feature-pipeline worker writes the output row for click event `e42`, then crashes before it records that the write succeeded. Its replacement cannot tell whether `e42` was committed. Skipping the event risks losing it. Replaying it with an append writes the feature row a second time, so that click counts twice in every training set built from the table. Writing the row under the stable key `(feature_version, e42)` turns the replay into an overwrite of the same row, so the uncertain event can be reprocessed safely.
>
> **Systems lesson**: Exactly-once belongs at the output commit boundary, whose record must bind the input position, transformation version, stable identifier, and commit state.

The two replays differ only in how the output is addressed. An append gives every attempt its own row, while a write under a stable key sends every attempt to the same row, so reprocessing uncertain work leaves the stored result unchanged (figure 4.9).

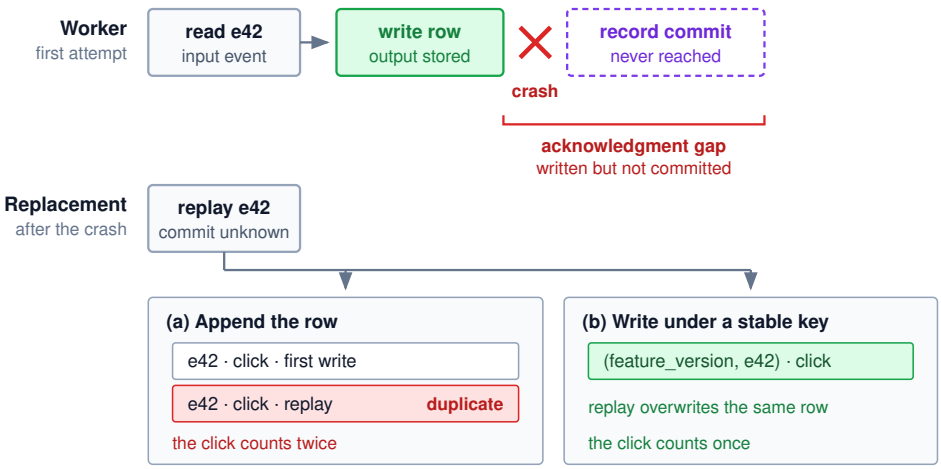

Figure 4.9: The Acknowledgment Gap at a Commit Boundary: A worker stores the output row for event e42 and crashes before it records the commit, so its replacement cannot tell whether the row exists and must replay the event. (a) Appending the replayed row duplicates it, and the click counts twice in every training set built from the table. (b) Writing under the stable key (feature_version, e42) makes the replay an overwrite of the same row, so at-least-once replay produces an effectively exactly-once result.

Streaming adds a time boundary to the contract. Event time records when an event occurred; processing time records when a worker handled it. A pipeline that aggregates clicks per hour must decide when an hour is finished, even though events can arrive late. A watermark is its running estimate that every event up to some event time has arrived, which lets it close that window and emit the result; a late-data policy determines whether later arrivals update, retract, quarantine, or drop the result (Akidau et al. 2015). Watermarks therefore trade freshness against completeness rather than prove that no older event remains.

Each of these mechanisms buys one property at the cost of another. Buffering absorbs finite bursts or temporary slowdowns but increases recovery state and freshness lag; sustained overload still requires backpressure or more service capacity. Aggressive watermarks reduce delay but exclude more late data, while replays help only when writes are safe to repeat. These choices change the stored result, because a feature table closed under an aggressive watermark holds different rows

from one closed under a lenient watermark. A reproducible training snapshot must therefore record source and code versions alongside partitioning, replay policy, stable identifiers, watermark state, and late-data decisions. That record is what lets a later feature backfill, a replay of history under new code, rejoin features at the correct points in time without mixing old and new histories.

↪ Feature backfills that replay history under this snapshot contract are developed in section 12.4.5.

4.7.2 The synthetic fuel line

The versioned-artifact contract applies to generated training samples as well as collected ones. When a fleet generates candidate samples, filters them, and verifies the survivors, each stage adds storage traffic and records that must remain linked to the accepted payload. The pressure to generate comes from the compute-optimal allocation of section 1.3, which grows the token budget with the parameter count. As high-quality human-generated text becomes harder to expand, the fleet encounters the **data wall**: the point where additional training samples must increasingly come from synthesis, simulation, curation, or expensive human collection (Villalobos et al. 2024).

This shift from simple collection to managed data generation sends more than payload bytes through the hierarchy. A "synthetic fuel line" must also store the **provenance chain** of every sample: which model generated it, which filters accepted or rejected it, which verifier models scored it, and which downstream run consumed it. That lineage makes repeated reuse of generated material traceable, and the data wall invites that reuse, because a fleet short of new human-generated text can end up training on its own earlier outputs. The lineage record therefore supports investigation of **model collapse**, where recursive training on generated data can degrade the learned distribution (Shumailov et al. 2024); retaining provenance alone does not prevent the degradation. The storage design must separately budget retained artifacts and the repeated scans used to assess them.

Napkin Math 4.8: The storage cost of synthetic data

Problem: A 1 TB synthetic dataset retains provenance metadata and undergoes three independent verification scans. How much persistent capacity and pipeline traffic does it require if verifier outputs are not retained as full dataset copies?

1. **Raw Payload**: 1 TB.
2. **Provenance Overhead**: 40 percent extra for lineage hashes, generation logs, and verifier scores.
3. **Verification Traffic**: 3 verifier passes each read the payload, producing 3 TB of read traffic. Their scores reside in the provenance metadata rather than in three retained payload copies.
4. **Capacity Result**: Persistent capacity is 1 TB $\times$ 1.4 = 1.4 TB, a 1.4× capacity factor.
5. **Traffic Result**: Writing the payload and metadata once and reading the payload three times moves 4.4 TB, a 4.4× traffic factor before generation and verifier compute are counted.

Systems insight: Verification primarily draws on service rate and compute, while retained provenance consumes capacity. The two costs coincide only when the pipeline materializes each verifier output as another full dataset. Synthetic storage must therefore budget bytes at rest separately from bytes moved and record which generator, filters, verifiers, and downstream runs produced each accepted sample.

Generated data therefore makes lineage part of the payload, so the store must budget retained provenance and verification traffic alongside the samples themselves.

4.8 Storage Economics

Checkpoints, training shards, and the lineage that binds them now each have a place in the hierarchy, and every placement carries a price set by the tier that holds the bytes and by how often they cross a billing boundary. ML storage economics is therefore not simply about choosing the cheapest capacity tier; the design objective is to minimize total cost of data delivery, which combines recurring storage

charges, cross-boundary transfer fees, request and retrieval fees, and the opportunity cost of idle accelerators.

Applied to one 100 TB training dataset, the per-gigabyte rates quoted earlier for each tier diverge by orders of magnitude: object storage (S3 Standard) costs $24,000 per year, provisioned local NVMe costs $120,000 per year, and archival storage (such as Glacier) drops to $4,800 per year. By contrast, hypothetically provisioning the same capacity as GPU memory at $15/GB would require $1,500,000 in upfront acquisition cost alone, before factoring in amortization, power, or server housing.

The 5× cost difference between local NVMe and object storage motivates placing data on the least expensive tier that meets its delivery and recovery requirements. Device technology contributes to the cost gradient. HBM uses 3D-stacked silicon with through-silicon vias, and accelerator-adjacent capacity is constrained by packaging and node design. Shared storage pools capacity across clients but introduces network and service dependencies. The cost comparison must include those dependencies rather than infer the cheapest usable path from capacity price alone.

Repeated cross-boundary egress, not storage, dominates cloud cost.

Storage cost, however, is only part of the equation. Data transfer costs can dominate, especially in cloud environments. When compute runs in a different region (a provider's separate group of data centers) or cloud from the bucket, as assumed here, reading the 100 TB dataset from object storage to compute instances incurs a transfer charge of $9,000, called an egress fee because it bills bytes leaving the region or cloud. For multi-epoch training that reads the dataset 10 times, the egress cost alone reaches $90,000, more than the annual storage cost. This inversion, where *reading* data costs more than *storing* it, drives the architecture decision to cache data on local NVMe rather than streaming from object storage each epoch.

The economic analysis changes further when accounting for the opportunity cost of accelerator idle time. The book's reference cluster of 1,024 GPUs, half the chapter's running cluster, at $2/GPU-hour costs $49,152 per day. If input stalls reduce its compute-active fraction from 90 percent to 70 percent, the lost compute time is valued at $9,830 per day. Over a 30 days training run, this amounts to $294,912. That value sets a budget for evaluating a storage upgrade, but the recoverable benefit depends on how much exposed wait the upgrade removes. Capacity that does not relieve the limiting stage yields no corresponding compute-time saving.

> Systems Perspective 4.3: The storage cost iceberg
>
> The visible cost of storage ($/GB/month) is the tip of the iceberg. Below the surface lie costs that often exceed the storage cost itself:
>
> - **Egress fees**: The ten-epoch egress charge of $90,000 computed above already exceeds a year of storage.
> - **Request fees**: Object storage charges per-request fees. S3 Standard GET requests cost about $0.0004 per 1,000 requests, so a dataset of 100M individual files, read once, costs about $40 in GET request fees before data transfer or retrieval charges, even though the storage cost is $2,000/month.
> - **Idle accelerator cost**: The $9,830 per day of lost compute computed above dwarfs any storage saving.
> - **Retrieval fees**: Archive storage charges retrieval fees ($0.02/GB for Glacier) in addition to egress. Restoring a 100 TB dataset from Glacier costs $2,000 in retrieval fees plus $9,000 in egress.
>
> The economically rational strategy often invests more in storage (local NVMe caching, parallel file system capacity) to reduce transfer costs and prevent accelerator stalls.

The cost iceberg reveals that raw storage capacity can account for less than half of the total expenditure; replication, metadata operations, API calls, and cross-region transfers can collectively dominate the bill at scale. Among those hidden line items, egress is the one that often flips the architecture decision. Unlike capacity charges, which scale with the volume of data stored, egress charges scale with the volume of data read, and ML training reads the full dataset on every epoch.

Napkin Math 4.9: The cost of egress

Problem: A team trains a vision model on a 50 TB image dataset stored in S3. Training runs for 20 epochs. The decision is whether to stream from S3 each epoch or stage to local NVMe.

Option A: Stream from S3

- Storage: 50 TB × \$0.02/GB/month × 12 months = \$12,000/year
- Egress: 50 TB × 20 epochs × \$0.09/GB = \$90,000 per training run
- Total for 4 training runs/year: \$12,000 + \$360,000 = \$372,000/year

Option B: Stage to local NVMe

- Storage (S3 source of truth): \$12,000/year
- NVMe capacity (50 TB across cluster): 50 TB × \$0.10/GB/month × 12 = \$60,000/year
- Egress (stage once per training run): 50 TB × 4 runs × \$0.09/GB = \$18,000/year
- Total: \$12,000 + \$60,000 + \$18,000 = \$90,000/year

Systems insight: Local NVMe caching saves \$282,000/year at 20 epochs per run, enough to fund 9 additional H100 GPUs at list price. Including the \$60,000/year provisioned NVMe cost and the once-per-run staging egress, the break-even point is just over 4 epochs per run; by the fifth epoch, local caching is cheaper than streaming from S3 every epoch.

The notebook turns the iceberg into a decision rule: the cheapest storage tier is not always the cheapest delivery path. Object storage can remain the source of truth, but repeated reads make egress, request overhead, and idle-accelerator cost part of the storage design. Multi-epoch vision training may justify staging on local NVMe by the fifth epoch, while single-epoch language-model pretraining may leave the corpus in object storage and optimize the streaming path instead. The running example stages anyway, although it reads its corpus only once, because egress is not what decides it. Each node's share of the shards fits easily on local NVMe, the staging copy happens once before the clock starts, and local reads keep object-store latency tails out of every training step. When compute and bucket share a region, transfer fees largely disappear and the same comparison turns on request charges, latency, and idle accelerators instead. The architecture is therefore chosen from access pattern, read count, and latency sensitivity rather than from \$/GB/month alone.

The cost analysis extends to the parallel file system tier, which occupies a middle ground between NVMe and object storage in both performance and price. A parallel file system capable of delivering 1 TB/s aggregate bandwidth requires roughly 100 OSS nodes, each with multiple NVMe drives and high-bandwidth network connections. Because the throughput target, not the stored capacity, sets that server count, a bandwidth-sized system costs far more than its stored bytes alone would suggest. This expense is justified only if the alternative, streaming from object storage and paying egress fees, is even more expensive, or if the latency requirements of checkpoint writes cannot be met by object storage alone. The break-even analysis depends on the organization's workload mix. A team running a single long training job may find that staging data on local NVMe and bypassing the parallel file system entirely is the cheapest option, while a team running many concurrent short jobs benefits from the shared namespace that a parallel file system provides.

The decision to build vs. buy storage infrastructure hinges on the trade-off between the high capital expenditure of an on-premises parallel file system and the recurring operational expenditure of cloud object storage. A 1-petabyte Lustre deployment, sized for capacity rather than for the 1 TB/s target above, costs roughly \$3–5 million in hardware, plus an additional \$500,000 per year in operational overhead for power, cooling, and engineering support. Storing 1 PB in a cloud object store at \$0.02/GB/month costs \$240,000 per year for storage alone. Egress fees for accessing the data, however, dominate the total cost. Reading that petabyte at 10 reads/year adds \$900,000 in data transfer fees.

Under these assumptions, annual cloud delivery costs \$1,140,000; after subtracting \$500,000 of on-premises operations, the net savings are \$640,000 per year. The resulting hardware break-even is roughly 4.7–7.8 years, and a payback within two years would require a hotter workload with many more full-dataset reads per year, lower capital expense, or non-egress drivers such as checkpoint

latency and shared-namespace requirements. For the running example, the 6 TB serialized corpus is much smaller than 1 PB, so the local-storage case is driven more by latency, repeated experiments, and checkpoint staging than by one-year egress savings alone.

4.8.1 Tiering strategies

The cost structure laid out above turns tiering into a placement policy: keep data only as close to the accelerator as its access frequency and latency requirement justify. A tiering strategy encodes that policy across the hierarchy. Its hot, warm, cold, and archive labels name placement roles for a dataset rather than devices, and they follow the scope brackets of figure 4.1. HBM stays outside these labels, because it holds the live batch and model state rather than a stored dataset.

The hot tier is the node-local pair of local NVMe and host DRAM, which holds the data the running job consumes. Earlier sections called local NVMe a warm cache because it sits in front of shared storage; in this placement vocabulary the same drives belong to the hot tier. This tier is provisioned per node, so data that multiple jobs or teams need simultaneously cannot live here alone. That shared-access requirement justifies the warm tier, the parallel file system, whose shared namespace and strong consistency serve datasets used by multiple jobs or teams, shared checkpoints for fault recovery, and model artifacts under active development. The cold tier is object storage, the canonical repository where datasets are authored and versioned and the durable backup for checkpoints staged through the parallel file system. Data that must be retained for compliance or reproducibility but is not expected to be accessed during normal operations drops to the archive tier (Glacier or equivalent).

Lifecycle policies of the kind shown for the archive tier in section 4.3.8 automate both directions of movement. They promote data from cold to warm when a job requests it, demote it when no job has read it for a configured period, and transition cold data to archive after 90 to 365 days since last access.

In the running example, the 3 TB source corpus and 6 TB serialized shards live permanently in object storage (cold tier), and the staging and cleanup lifecycle of section 4.5.1 moves each node's shards through local NVMe (hot tier) for the life of the job. The parallel file system holds only the job's recent checkpoints, which the retention policy of section 4.6.2 governs.

The economic benefit of automated tiering compounds over time. An organization retaining 50 training datasets of 10 TB each would need 500 TB of parallel file system capacity if all data were left in place. If ten datasets are active concurrently and the other forty can be demoted after their jobs finish, automated tiering reduces the dataset capacity required on the parallel file system to 100 TB, saving $144,000 per year at the parallel file system's per-TB rate. These savings are invisible in any single experiment but substantial when accumulated across the fleet. Tiering and staging thus settle where training data and checkpoints live and what each placement costs; serving reuses the same hierarchy for a different objective, governed by the time to bring a model into HBM and the latency of serving-time state rather than epoch throughput.

4.9 Storage for Serving

Where full-corpus training streams datasets for weeks and writes checkpoints in bursts, the weight-resident serving baseline of section 4.2 reads the model weights once when a replica starts or recovers and then answers requests out of HBM, before KV-cache offload and retrieval indexes add serving-time I/O of their own.

That startup read makes model loading latency, often called cold-start time, the primary storage metric for weight-resident inference: the duration required to load a model from storage into the accelerator's HBM. For the 175B parameter language model, the weights alone occupy 350 GB in FP16 format; loading that footprint sequentially from a single high-performance NVMe drive at 14 GB/s takes 25 seconds, an unacceptable delay for a user-facing application. Loading from a shared PFS at a per-node rate of 4 GB/s takes 87.5 seconds. Reading from object storage at 1 GB/s stretches the delay to 5.8 minutes.

To reduce cold-start time, the model is sharded and loaded in parallel. If the 350 GB model is striped across 8 NVMe drives, each loading a 43.75 GB shard at 7 GB/s, the total load time drops to roughly 6.2 seconds. When using tensor parallelism across 8 GPUs, each GPU loads only its 43.75 GB shard. If the weights already sit in host DRAM, each GPU's own PCIe link could move its shard

in about 0.7 seconds, but the 8 transfers share the node's 200 GB/s of host-memory bandwidth, so the full load takes about 1.8 seconds. For the most latency-sensitive applications, organizations maintain warm replicas: one or more copies of the model weights kept preloaded in host DRAM, ready for that host-memory-bound transfer. This trades DRAM capacity for a load of about 1.8 seconds instead of the 6.2 seconds striped NVMe read.

These load times describe one replica. The model rollout previewed in section 4.2 asks all 100 replicas to load a new version at once, 35 TB of weights in one burst. If every replica reads from object storage at the 1 GB/s single-client rate, the fleet asks the store for 100 GB/s. A store whose aggregate service matches that figure, the 100 GB/s of table 4.2, finishes the rollout in about 5.8 minutes, no slower than one replica's cold start; any smaller aggregate share stretches the rollout in proportion. Reading from the 1 TB/s parallel file system instead would take about 35 seconds, but only by claiming the whole file system and starving every training job that shares it, the read-side mirror of a checkpoint storm. Serving fleets therefore stage the new version on local NVMe at a paced rate ahead of the traffic switch, roll the update across groups of replicas so that the old version keeps serving, or let replicas that have finished loading supply the rest.

↪ How cold-start time limits autoscaling, and how warm pools trade idle capacity for scaling response, is developed in section 10.13.

During autoregressive generation, the server keeps a per-request record of the attention keys and values produced by earlier tokens so it can generate the next token without recomputing the entire prefix. This record, called the key-value cache (KV cache), grows with sequence length and can consume significant HBM. For long sequences (128K tokens or more), KV cache offloading to host DRAM or NVMe extends the effective context window at the cost of increased latency per generated token, because cache blocks must be moved back into HBM before attention can use them. The storage lesson is that inference maintains an active working set rather than presenting solely a model-loading demand. When serving must also retrieve external context to ground that generation, the storage pipeline faces an entirely different access pattern: low-latency random queries into high-dimensional vector indexes.

↪ KV-cache growth and the HBM capacity wall it creates for serving are developed in section 10.7.1.

4.9.1 Retrieval indexes

An inference request that needs supporting documents cannot wait for a sequential scan of the training corpus. A retrieval system first converts documents into vectors that represent semantic meaning, then finds nearby vectors and returns their source documents as context. The same storage analysis still applies, but the unit of demand becomes a query and its index accesses, and the deadline becomes the serving latency budget. With a graph-based approximate nearest-neighbor (ANN) index, which links each vector to a few of its neighbors, the storage system traverses a high-dimensional vector graph instead of streaming sequential blobs via formats such as TFRecord or Parquet. Other index families shift the storage demand elsewhere. Exact flat search scans every stored vector and is bound by scan bandwidth. Inverted-file (IVF) partitioning groups vectors into clusters and probes only the few nearest the query, so each query reads a small part of the index. Product quantization (PQ) compresses each vector into short codes so that more of the index fits in memory, at some cost in accuracy.

> Systems Perspective 4.4: The IOPS bottleneck in high-dimensional search
>
> For graph-based ANN, this architectural pattern shifts the primary storage bottleneck from sequential read bandwidth to random-access IOPS. The search traverses a high-dimensional graph to find approximate nearest neighbors across billions of embeddings. Its performance is governed by the trade-off among recall, latency, and capacity, where recall is the fraction of the true nearest neighbors that the approximate search returns. Raising recall requires traversing a larger portion of the index, driving up IOPS and latency. The arithmetic of a disk-resident traversal shows why IOPS, not bytes, sets the limit. Each hop that misses DRAM costs one dependent SSD read of roughly 100 µs (table 4.2), and the next hop cannot start until that read returns, so ten such hops already spend a millisecond of the serving budget in storage alone. Throughput is bounded separately. A drive whose 4 KB random reads run near 0.5 GB/s (section 4.2) completes on the order of 10^5 such reads per second, and every concurrent query draws on that same budget, while each read moves only a few kilobytes.

How often a hop misses DRAM depends on how much of the index fits there. A 1,536-dimension embedding in FP32 (4 bytes per value) consumes approximately 6.1 KB, so a billion-document corpus requires roughly 6.1 TB before graph adjacency lists. At 50 to 100 bytes per node, those lists add 50 GB to 100 GB, for about 6.2 TB at the high end. Production indexes can be larger once identifiers, graph levels, allocator overhead, deleted-entry slack, and replication are included. In-memory hierarchical navigable small world graph indexes make graph navigation fast when the index fits in DRAM (Malkov and Yashunin 2020). Disk-backed graph indexes such as DiskANN trade DRAM footprint against SSD traversal while targeting high recall and low query latency (Subramanya et al. 2019).

Graph-based retrieval therefore spends storage budget on index levels that keep random traversal inside the serving latency service-level objective. A disk-backed graph is economical only if the extra SSD accesses do not erase the quality gain from searching a larger corpus. The same hierarchy now supports bulk model loading, request-scoped KV state, and retrieval, each with a distinct service-rate contract. Every design so far also assumes today's six tiers, and the boundaries between those tiers are themselves moving.

4.10 Storage Technologies That Move Tiers

Today's six-tier hierarchy leaves a 1,000× latency cliff between host DRAM (~100 ns) and local NVMe (~100 µs), forces the host CPU to spend cycles decoding and staging payload bytes, and pauses training while synchronous checkpoints reach flash. Compute Express Link (CXL), the coherent capacity tier of section 2.9.1, targets the DRAM-to-NVMe gap first. Seen from storage, CXL-attached memory is a "Tier 1.5," slower than host DRAM but byte-addressable, with access latencies of roughly 200 to 500 ns against the roughly 100 µs of an NVMe read. It offers a capacity pool for the structures that the Tier 1 discussion showed competing for host DRAM. For the running example, CXL-attached memory could expand the effective host memory tier from 512 GB to over 4 TB per node. That pool could hold deeper prefetch queues, embedding-table spillover, or a local shard cache as large as the 3 TB compressed source corpus without paying NVMe access latency on every read, though it still falls short of the full 6 TB tokenized training copy on one node.

A different bottleneck appears when CPU-side decoding, not storage bandwidth, throttles the pipeline. The computational storage drive relieves that bottleneck by embedding processing elements, either field-programmable gate arrays (FPGAs) or simple CPU cores, directly on the SSD controller, enabling decompression, filtering, or format conversion to happen at the storage device rather than consuming host CPU cycles. For ML pipelines where CPU-side decoding is the bottleneck, computational storage could eliminate the CPU from the data path entirely, complementing GDS by offloading work that direct DMA cannot address.

Checkpoint writes impose the pause T_{write} while a consistent snapshot reaches the first storage tier, and shrinking it makes frequent saves less costly. Persistent memory, despite market shifts such as Intel's 2022 wind-down of its Optane memory business (Intel Corporation 2022), remains the architectural idea that attacks that term directly. Byte-addressable non-volatile memory could lower the destination-device latency, but the checkpoint pause still includes snapshot creation, byte transfer, and coordination. A lower device-access latency reduces T_{write} only to the extent that this device lies on the pause's critical path.

These technologies are not yet broadly deployed at the scale of the largest ML clusters, but they illustrate the direction of the storage hierarchy's evolution. The fundamental principle remains unchanged, with faster, more expensive storage closer to the accelerator and slower, cheaper storage at the periphery. What changes is the *granularity* of the tiers and the *size of the gaps* between them.

4.11 Fallacies and Pitfalls

Storage design errors are among the most expensive mistakes in ML infrastructure because they are discovered late, when training begins, and remediation requires either expensive hardware upgrades or time-consuming data reformatting. The following fallacies and pitfalls capture the most common errors that experienced storage engineers make when first encountering ML workloads.

Fallacy: *NVMe is fast enough to feed GPUs directly without pipelining.*

A device can have enough sustained bandwidth and still expose stalls when transfers are serialized with computation. Without pipelining and prefetching to hide the latency of loading from NVMe to host DRAM to HBM, every batch transfer introduces a stall equal to the transfer time. Comparing a four-drive RAID-0 stripe with HBM, a gap of about 120×, does not show whether NVMe can feed the job, because model state is reused many times for each batch loaded from storage. The relevant test compares measured input delivery with batch bytes divided by compute time, then checks whether buffering hides transfer latency and jitter.

Pitfall: *Relying on prefetching to hide object storage latency.*

Prefetching hides latency only within the time covered by ready batches; it cannot hide an arbitrarily long interruption. With N_{prefetch} batches ready and constant compute time, a complete interruption of supply exhausts the queue after $N_{\text{prefetch}} \times T_{\text{compute}}$. A slow request need not stall computation if other completed work remains available, but correlated delays can deplete many workers' queues together. Provision buffer depth from measured delivery tails, test correlated interruptions, and verify that the sustained refill rate meets consumption. Local staging can reduce remote exposure when the required shards fit and remain valid.

Fallacy: *Millions of small training files are harmless at scale.*

Storing each training sample as an individual file (one JPEG per image, one JSON per text sample) creates repeated per-file work through operations such as `open()`, `stat()`, and `close()`. Metadata becomes the bottleneck when the resulting request rate exceeds the service's capacity, even if byte bandwidth remains available. Aggregating samples into sequential shards amortizes that work (section 4.3.5). Packaging 10,000 to 100,000 samples per shard can reduce file-level operations by a corresponding 10,000 to 100,000× when readers consume whole shards. The loader must still implement the intended sampling policy and avoid repeated enumeration.

Pitfall: *Ignoring egress costs when designing the storage architecture.*

Cloud egress costs are invisible during development but can dominate at scale. Under a representative $0.09/GB cross-boundary egress fee, a training job that reads a 100 TB dataset from object storage 10 times incurs $90,000 in egress fees, far exceeding the annual storage cost of $24,000 under the object-storage price assumption in section 4.3.7. These fees apply when reads cross a region, provider, or internet boundary; same-region reads may be free or discounted, which is exactly why the cost stays hidden until a cross-region or multi-cloud topology makes it unavoidable. Teams that prototype with small datasets on object storage and then scale up discover that their storage architecture is economically unsustainable. The fix is to budget egress costs explicitly and to design the data pipeline to minimize cross-tier transfers, typically by staging data on local NVMe at job start rather than streaming from object storage each epoch.

Fallacy: *Average write bandwidth is enough for checkpoint design.*

Checkpoints are written by all nodes simultaneously, creating bursts that can exceed steady-state bandwidth by 10× or more. When the PFS write lies on the training critical path, sizing it only for average write load extends T_{write} during bursts and reduces effective training throughput. A synchronous PFS-direct design must provision enough burst bandwidth to meet its pause budget. With local staging, the first tier absorbs the burst and the PFS drains it in the background; the shared tier must then meet the durability deadline and keep up with repeated saves without an unbounded backlog.

Pitfall: *Assuming faster accelerators automatically improve training throughput.*

When an organization upgrades from A100 to H100 GPUs, they expect training time to decrease proportionally to the compute improvement. If the storage pipeline was already marginal on A100s, delivering data just fast enough to avoid stalls, the faster H100 compute hits the I/O wall sooner. The accelerator upgrade reduces the compute time T_{compute} without reducing the input-pipeline time T_{IO}, which means the data stall ratio increases. Every accelerator upgrade must be accompanied by a storage pipeline audit to ensure that the faster compute does not expose a previously hidden storage bottleneck.

Fallacy: *RAID configurations optimized for databases are suitable for ML training.*

RAID selection follows the tier's failure contract. A disposable local cache can use RAID 0 to aggregate drives without parity, provided a durable source remains available and the fleet can tolerate restaging after a drive failure. Redundant layouts can still reduce disruption when restaging is slow or the cache supports an availability objective. A local checkpoint awaiting durable replication is not yet disposable merely because it lives on the same NVMe drives. Treating all ML storage as replaceable input data confuses capacity placement with recovery policy.

Pitfall: *Creating full copies for each dataset snapshot.*

Copying an entire dataset for every version duplicates unchanged content. For a 10 TB dataset with 5 percent of its bytes changed, a full copy retains 9.5 TB of redundant data. Content-addressable storage or delta encoding can reuse unchanged objects or blocks. If the storage granularity isolates the changed bytes, the new payload is roughly 500 GB, a 20× reduction before metadata and redundancy overhead. Rewriting large compressed shards can reduce that saving because a small sample change may replace a much larger stored object. Versioning must therefore preserve lineage while choosing a reuse granularity compatible with the dataset format.

Fallacy: *Compression always helps because it reduces the amount of data to read.*

Compression reduces stored and transferred bytes but adds decoding work. If decoding further saturates an already limiting host processor, it makes that bottleneck *worse*, not *better*. If transfer is limiting and decoding has headroom, compression can improve delivery until another stage becomes the limit. Compare candidate representations using output-equivalent stage rates, including CPU or GPU decoding, rather than select the strongest codec from an I/O-bound label alone. The same choice can reduce capacity cost without improving training throughput.

These fallacies and pitfalls all return to the same accounting exercise: where each byte lives, how often it moves, and which tier pays the cost.

4.12 Summary

Storage in ML systems is an active, multi-tiered pipeline that supplies accelerator inputs, preserves restorable state, and records the lineage of training and serving artifacts. The hierarchy exists because workloads need different combinations of bandwidth, capacity, latency, durability, and cost, and the demand on it grows with the scaling laws that drive the fleet. Larger models call for more training tokens and carry larger checkpoints, so storage pressure rises at every step up the model-size curve. The data pipeline throughput equation, $\text{BW}_{\text{required}} = N_{\text{GPU}} \times \eta_{\text{target}} \times D_{\text{vol}}^{(\text{batch})} / T_{\text{compute}}$, sizes delivery for a chosen training rate when applied to the representation and distinct input consumers at each tier. Measured against the Fleet Law, storage succeeds when it contributes no exposed time. Prefetching preserves the Fleet Law's assumption that the next batch is already resident. Tiered checkpoint staging shrinks the save pause that surfaces in $T_{\text{sync}}(N)$ to the local write and takes the slow shared-storage copy off the step entirely, so that copy costs step time only when it contends with collectives or input reads. Both spend communication (C_2) and coordination (C_3) capacity to protect compute (C_1).

Treating the hierarchy as a cache makes that sizing workload-specific. The hit rate and the service times of the hot and cold paths determine mean effective throughput, while the tail and correlation of misses determine whether a finite prefetch buffer drains. At the durable tiers, reliability also separates into three contracts: data survival, service response, and the version a reader may observe. These are durability, availability, and consistency, respectively; improving one does not automatically establish the others.

The diagnostic task is to identify which service contract fails before changing the storage tier. Low accelerator utilization can arise from metadata-intensive loading, a depleted prefetch queue, an exposed checkpoint write, or host-memory traffic crossing a contended NUMA link. Measure input wait separately from communication and synchronization before attributing it to storage. Full-corpus scans favor aggregated shards, while graph retrieval and live features require random access and explicit consistency. A depleted input queue can reflect slow storage, metadata contention, or insufficient decoding capacity. For checkpoint writes, evaluate the shorter training pause together with the time until a complete generation becomes durable and publishable.

Key Takeaways: Feed the accelerators or waste them

- **Workload demand sizes the delivery path**: Use $\text{BW}_{\text{required}} = N_{\text{GPU}} \times \eta_{\text{target}} \times D_{\text{vol}}^{(\text{batch})} / T_{\text{compute}}$ for the chosen training rate, then check every stage's service rate and buffering, because delivery is the minimum over stages and faster compute against a fixed input path plateaus at that rate (principle 7). GDS removes the host payload copy, not CPU control of the I/O, and helps when that copy, or the CPU time it takes from other stages, limits delivery.
- **Workload determines the storage contract**: Full-corpus scans favor sequential throughput and aggregated shards, and metadata operations can limit delivery before byte bandwidth does. Disk-backed graph retrieval must fit its random I/O within a serving deadline, and synthetic-data pipelines must budget verification traffic separately from retained capacity while keeping the versioned lineage of every accepted sample.
- **A tier ratio is not a feed requirement**: Tiers quote bandwidth per accelerator, per node, or as a shared aggregate, so divide shared bandwidth among active clients and size delivery from input bytes rather than from the gap between HBM and storage (principle 8).
- **Tail latency sets the cache budget**: A high hit rate is not enough when misses are orders of magnitude slower. Served serially, a 95 percent hit rate with misses 100× slower yields only about one-sixth of the all-hot rate. Use equation 4.1 as a mean baseline, then size prefetch buffers from p99 I/O latency and test correlated miss bursts against the compute window.
- **Checkpoint recovery is a storage protocol**: Local NVMe staging can shorten pause time, but a restorable checkpoint also needs durable publication, a committed manifest, and verified shards; durability, availability, and consistency remain distinct contracts.
- **Pipeline correctness lives at commit boundaries**: Replay preserves processing completeness only when stable identifiers and idempotent or transactional writes prevent duplicate effects. Event-time watermarks then make the separate freshness-completeness trade-off explicit.
- **Total delivery cost, not price per byte, decides tiering**: Capacity charges scale with bytes stored while transfer charges scale with bytes read, so a dataset read many times across a billing boundary can cost more to read than to store (principle 8). Compare total delivery cost, including transfer charges and idle accelerators, before deciding where to stage data.

Checkpoint staging shows the C^3 trade-off in one design decision. Writing first to local NVMe shortens the compute pause, but asynchronous replication adds communication traffic and requires coordination to publish only complete generations. Once the shared-file-system write leaves the critical path, the bottleneck can move to local flash endurance, peer-replication bandwidth, or manifest commit latency. Measuring only the shorter training pause would miss the new recovery boundary.

What's Next: From storage to communication

Compute silicon, network fabrics, and storage hierarchies establish the physical resources of the machine learning fleet. Coordinating those resources into one training computation remains a separate problem. This chapter's 175B running model already exceeds, in its FP16 weights alone, the 80 GB device of the HBM placement example, so its state and work must span devices, and every division of that work obliges workers to exchange gradients, activations, or parameters across the fabric. Part II moves from physical infrastructure to distributed execution. It begins with the collective operations that carry those exchanges and the cost of each, then

partitions neural networks across data, tensor, pipeline, and sequence dimensions, before turning to the hardware failures that a fleet must survive.

↪ The collective operations that carry gradient and activation exchanges, and the cost models that price them, are developed in chapter 5.

II

DISTRIBUTED ML

Part II

4

Distributed ML Principles

Wiring thousands of accelerators together with high-speed switches and parallel storage does not by itself make them train one model. Without distributed coordination, replicas diverge, partitioned layers stall on missing activations, and a single failed GPU halts the entire cluster. If Part I built the physical substrate, Part II establishes the logic of distribution: the algorithms, protocols, and coordination strategies that make a collection of independent accelerators behave as one machine. The same physics that set the rates in Part I explains why adding accelerators does not always yield linear speedups, and why communication, not computation, is often the hidden bottleneck.

At this scale, the engineering challenge is a trade-off between parallelization and coordination. Partitioning a model reduces per-device memory pressure but requires exchanging intermediate results across partitions. Replicating a model across workers instead requires coordinating their updates. Adding nodes can reduce training time while exposing a job to more component failures, decreasing its mean time between failures (MTBF) when an individual failure interrupts the job. Fault tolerance therefore becomes part of the execution design. These principles decompose the communication cost that the Fleet Law exposes (section 1.5.4) and add the reliability cost of scale.

The following principles preview the trade-offs developed in Part II, and the chapters establish when each mechanism helps and how to measure its cost. The first pressure is the economics of scale itself. Frontier capability demands compute, and each improvement costs more than the last, as the scaling laws of section 1.3 showed. Part I priced the supply of that compute; the first principle states the demand in the form that later chapters cite.

Principle 9: The empirical scaling law

Invariant: Within empirically validated regimes, foundation-model loss ($\mathcal{L}$) follows approximate power-law relationships with a resource quantity X such as training compute, dataset size, or parameter count. The exponent α_{scale} depends on the fitted regime. The single-resource simplification used to illustrate diminishing returns is:

$$\mathcal{L}(X) \propto X^{-\alpha_{\text{scale}}}$$

Implication: Within a fixed scaling regime, reducing loss further requires increasing resources. The diminishing return helps explain the move from single-server training to warehouse-scale clusters. Architecture, data quality, and training methods can change that regime, so a fitted scaling relation guides resource planning without guaranteeing improvements on every task.

That demand for scale turns each training step into a negotiation between faster local work and slower global coordination. Splitting a batch or a model across more workers shrinks the arithmetic each device performs, but it also introduces network transfers and barrier waits that grow with cluster size.

Principle 10: The fleet law (distributed step-time model)

Invariant: Under ideal partitioning of a fixed amount of work across N workers or accelerators, the step-time model adds computation, communication, and synchronization, then subtracts the portion hidden by overlap, with $0 \leq T_{\text{overlap}} \leq \min\left(\frac{T_{\text{compute}}}{N}, T_{\text{comm}}(N) + T_{\text{sync}}(N)\right)$.

$$T_{\text{step}}(N) = \frac{T_{\text{compute}}}{N} + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}}$$

Implication: Scaling is a race between parallelizable compute (which shrinks with N under ideal partitioning), communication overhead, and synchronization overhead. To scale efficiently, algorithms must reduce or amortize $T_{\text{comm}}(N)$ and $T_{\text{sync}}(N)$ (for example, through gradient accumulation, which alters the effective batch and convergence regime, or through compression, which can introduce numerical error). They must also maximize T_{overlap} through communication hiding and pipelined execution.

For data parallelism, the AllReduce that combines every worker's gradients sets the size of the Fleet Law's communication term. Because every replica must receive the full sum of gradients before advancing its weights, the volume moved per rank approaches a constant floor even as more workers join the job.

Principle 11: The AllReduce bandwidth bound

Invariant: An AllReduce of a M-byte tensor across N ranks, built from point-to-point transfers, must send at least $2(N-1)M/N$ bytes from some rank. Ring AllReduce meets that bound, at the price of a startup term that grows linearly with N:

$$T_{\text{ring}} = 2(N-1)\alpha + 2\frac{N-1}{N}\frac{M}{\beta}$$

Implication: The bandwidth term approaches $2M/\beta$ and stops shrinking as workers are added, while each worker's share of the compute keeps shrinking, so under pure data parallelism communication eventually dominates the step. Tree and hierarchical schedules give up some of the ring's byte efficiency for fewer startup rounds when messages are small or N is large.

The Fleet Law and the AllReduce bound predict the time of one step; the scaling efficiency bound measures how far the resulting speedup falls short of ideal. Comparing multi-device runtime against a single-device baseline turns those exposed communication and synchronization terms into a direct measure of cluster return on investment.

Principle 12: The scaling efficiency bound

Invariant: For a fixed amount of training work, adding nodes reduces the computation assigned to each node, but the communication and synchronization terms of the Fleet Law limit the resulting speedup.

$$\eta_{\text{scaling}} = \frac{T_1}{N \times T_N}$$

Here, T_1 and T_N are the times to complete the same work on one and N devices, respectively. The ratio measures achieved speedup as a fraction of ideal linear speedup.

Implication: Perfect linear scaling ($\eta_{\text{scaling}} = 1.0$) is the ideal baseline, not a practical target. Real systems fall short of it by an amount that depends on the workload, the fabric, and the parallelism layout. Under pure data parallelism the shortfall grows with N because the compute term keeps shrinking while the ring's bandwidth term stays near $2M/\beta$, the scaling

wall of data parallelism. A statistical limit applies as well, because beyond the critical batch size added data-parallel width buys little further loss reduction. The gap between $\eta_{\text{scaling}} = 1.0$ and the achieved efficiency captures losses from communication and coordination as well as load imbalance or other limits on parallel execution.

Even when communication and coordination are tuned, a larger fleet fails often enough that progress itself must be protected. Without periodic state snapshots, a single device or link failure discards every step completed since the run began.

Principle 13: The Young-Daly checkpoint interval

Invariant: The first-order checkpoint model balances the cost of writing the checkpoint (T_{write}) against the expected cost of reworking lost progress due to system-level failures ($\text{MTBF}_{\text{system}}$). For long runs with approximately constant failure rates and checkpoint write times small relative to MTBF, its optimal checkpoint interval (T^*_{ckpt}) is:

$$T^*_{\text{ckpt}} = \sqrt{2 \cdot T_{\text{write}} \cdot \text{MTBF}_{\text{system}}}$$

Implication: Checkpointing consumes time and storage bandwidth. When more participating components reduce job MTBF, the model favors shorter intervals, and the overhead at the optimum grows roughly as the square root of the component count; faster checkpoint writes can reduce that overhead. Correlated failures, variable write times, and recovery requirements need additional analysis before the interval becomes an operational schedule.

Fault tolerance illustrates why reducing one cost requires checking the others. Saving more often reduces expected rework but increases storage traffic.

Principle 14: The displacement of overhead

Invariant: Distributed ML optimizations often relocate overhead among Compute, Communication, and Coordination (the C^3 taxonomy). Reducing one term can increase another, so designs must evaluate all three.

Implication: Asynchronous training can eliminate global coordination barriers ($T_{\text{sync}}(N) \to 0$) but may introduce gradient staleness that shows up as additional training iterations ($T_{\text{compute}} \uparrow$). Pipeline parallelism can reduce communication volume relative to finer-grained model partitioning but adds pipeline bubble time ($T_{\text{sync}}(N) \uparrow$). The C^3 taxonomy reveals where each design pays its overhead.

Accounting for overhead inside one job leaves a fleet-level question open, because many jobs compete for the same machines and a share of the fleet is worth something only if the job can run on it. Granting accelerators without the interconnect topology, power, or storage bandwidth a synchronous workload requires only strands those devices in an unusable allocation.

Principle 15: The feasibility-before-fairness principle

Invariant: A gang allocation for a synchronous job requires atomic allocation (every member is granted at once) and topology-feasible placement (the members sit within the interconnect domains the job's parallelism requires) to complete training steps. Its power and I/O allocation must also support the rate of progress promised at admission. Missing required gang members or placement can prevent execution; insufficient power or I/O capacity can reduce progress below the admitted rate.

Implication: Fairness policies such as dominant-resource fairness divide useful progress, so they can choose only among feasible allocations. A scheduler must first enumerate the feasible bundles and only then apply its fairness policy to them; a fair share of accelerators on which the job cannot run is idle capacity charged to a tenant. Feasibility constraints can deliberately leave accelerators idle, and that idle capacity is the price of keeping every admitted job productive.

Part II develops these previews in dependency order. Collective communication establishes how groups of workers combine and redistribute information across the fabric and what each exchange costs. Distributed training then partitions a workload and decides which of those exchanges its partitions must make. Fault tolerance protects the resulting progress against failure, and fleet orchestration allocates shared resources to jobs with those communication and recovery needs, admitting only allocations the jobs can actually use. Orchestration decides how much of the fleet a job actually holds, and fault tolerance decides how much of that held time becomes retained progress. In Fleet Law terms, collective communication shrinks $T_{\text{comm}}(N)$ and enlarges T_{overlap}. Distributed training decides how T_{compute}/N is divided and what $T_{\text{comm}}(N)$ that division incurs. Fault tolerance and orchestration bound the synchronization and recovery time that no partitioning can remove. The same chapters also set the fleet-level envelope for serving, since replicas draw on the same scheduler, share the same failure rates, and need the same recovery machinery. Together, they develop the algorithmic and operational discipline that turns the physical fleet of Part I into a working distributed system.

5

Collective Communication

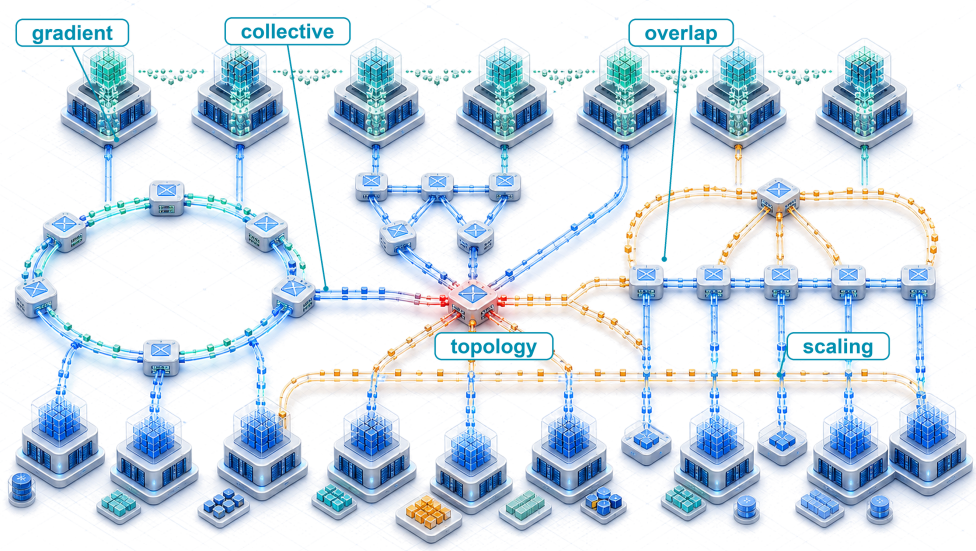

5.1 From Parallelism to Communication Patterns
5.2 The Physics of Data Movement
5.3 Parallelism Strategies and Their Traffic
5.4 Collective Operation Primitives
5.5 Pricing Collectives by Rounds and Chunks
5.6 AllReduce Algorithms and Topologies
5.7 Hierarchical Communication
5.8 AllToAll Exchanges at Scale
5.9 Gradient Compression Under Bandwidth Scarcity
5.10 Communication-Computation Overlap
5.11 Communication Libraries and Runtimes
5.12 Fallacies and Pitfalls
5.13 Summary

Purpose

Why does communication between machines constrain large-scale machine learning systems?

Computation scales by adding processors; communication scales by moving data between them. These scale differently. Adding processors increases available arithmetic capacity, but useful work stalls when the required data does not arrive in time. Exchanging gradients, activations, and parameters can then consume more time than computing them. On one machine, shortening the dominant calculation often shortens the job. Across a fleet, that same improvement may simply make each processor reach its next dependency sooner and wait longer for the others. The cost of that wait includes every reserved accelerator, so a small delay repeated throughout training can consume a substantial budget without advancing the model. Faster links help when transfer capacity is the constraint, but the benefit depends on which physical paths the traffic shares. More bandwidth cannot remove a dependency. Sending less data introduces its own choices about numerical accuracy, memory use, and when the next computation can begin. An exchange must also deliver the correct result to every participant that needs it. A fast transfer that combines the wrong contributions does not preserve a valid training run. The engineering decision is therefore how to organize communication so that its cost remains compatible with the work it enables. Light-speed delays, bandwidth limits, and the energy cost of movement make that decision as physical as processor design. In C^3 terms, collective communication connects distributed compute through coordinated exchanges whose correctness and exposed cost determine whether the fleet delivers useful scale.

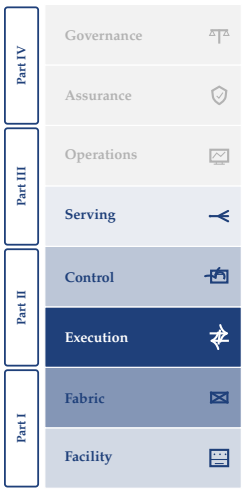

Learning Objectives

- Apply communication cost models to estimate collective duration and message-size crossover points
- Match collective communication primitives and AllToAll schedules to parallelism strategies, topology, message granularity, and skew
- Compare ring, tree, butterfly, double-tree, and hierarchical reduction algorithms using scale-dependent cost models
- Evaluate topology-aware collective schedules under link-bandwidth, rank-placement, and shared-capacity constraints
- Evaluate compression, sparsification, and error feedback by balancing communication savings against convergence risk
- Design overlap strategies with bucket fusion, asynchronous operations, and layer-by-layer gradient scheduling
- Specify matched measurements for calibrating collective cost models across algorithms, process groups, and topologies

5.1 From Parallelism to Communication Patterns

When 10,000 GPUs cooperate on one training update, the splits created by data, tensor, pipeline, and expert parallelism determine which results must cross device boundaries. Workers execute collective operations within participating groups and point-to-point transfers between pipeline stages. In the multi-tier fleet stack of figure 1.4, these exchanges occupy the Distributed Execution row of the Distribution band, where logical parallelism becomes physical traffic.

Every such split assumes that replicas, shards, stages, and experts can exchange gradients, activations, parameters, and tokens quickly enough for the optimizer to see one coherent training run. In data parallelism, each worker holds a full copy of the model and processes its own slice of the batch, so the copies must agree on one combined gradient before every update. Tensor, pipeline, and expert parallelism divide the model itself, so partial results must cross device boundaries within each step. Each of these exchanges is charged to the Fleet Law's communication term $T_{\text{comm}}(N)$, and the part of it not hidden behind computation counts against program goodput, the ratio of a step's ideal execution time to its actual time (section 9.6.1). Every way of splitting the model relocates work onto the network, and this chapter turns those exchanges into schedules whose latency, bandwidth, and topology costs can be computed.

↪ The partitioning strategies themselves, and the 3D parallelism cube that composes data, pipeline, and tensor parallelism into one configuration space, are developed in section 6.2.

The gap between that assumption and the wire reveals an asymmetry in how computation and communication scale. Computation is local because each GPU works on its own data, so aggregate compute grows with the number of GPUs. Communication is global because keeping the fleet synchronized requires information to cross physical links with finite latency, bandwidth, interconnect topology, and energy. Gradient synchronization is the central case, because every synchronous data-parallel step repeats it on the full gradient. Scaling laws (principle 9) make that case harder over time, because the capability they reward comes from more parameters trained on more accelerators, so the gradient to be reduced and the group that must agree on it grow together.

Managing that global traffic begins by separating an operation's result from the schedule that produces it. Section 3.1 named AllReduce, AllGather, and AllToAll as the traffic the fabric must carry; here they become contracts on results. AllReduce[1] combines corresponding tensor elements across participating workers, called ranks after their index in the group, and returns the reduced tensor to every worker. AllGather instead concatenates their shards, ReduceScatter combines contributions but leaves each worker with only a shard of the result, and AllToAll sends different data to different destinations. The parallel decomposition determines which contract each exchange needs; ring, tree, and hierarchical algorithms are alternative schedules for carrying that contract out.

For the initial gradient budget, consider a ring schedule: each worker passes tensor chunks to a neighbor, first accumulating their sums and then circulating the completed chunks. Each phase moves nearly one tensor's worth of data per worker, so the bandwidth budget is approximately

[1] **AllReduce**: The name comes from MPI and describes the result, not the schedule: a Reduce (combining every rank's data at one root) followed by a Broadcast (returning the combined value to all ranks). A direct gather to the root concentrates raw input traffic there, but a tree can aggregate contributions along the way. The algorithms in section 5.6.3 distribute the reduction without a root.

twice the gradient size divided by link bandwidth. Section 5.6.3 derives the detailed round count once the transfer model is established.

Definition 5.1: Gradient synchronization

Gradient synchronization is the collective communication step in synchronous data-parallel training in which every worker contributes its locally computed gradient tensor to an aggregate reduction, then receives the same reduced result so all model replicas apply an identical update.

1. **Significance**: Communicating all gradients of a 70B-parameter model in BF16 gives a 140 GB logical payload per replica per step. As a communication-only baseline, suppose every rank reduces that complete buffer. A ring across 1,000 GPUs at 50 GB/s per link has a bandwidth budget of approximately $2 \times 140/50 \approx 5.6$ s of transfer per step. Model partitioning shrinks each rank's share of that buffer, and only the portion of $T_{\text{comm}}(N)$ left exposed after overlap adds to step time in the Fleet Law (equation 1.3).
2. **Distinction**: Unlike a parameter server, which concentrates every worker's gradient at a central aggregator, ring AllReduce spreads the exchange so that each worker sends $2 \times (N-1)/N \times M$ bytes per step, approaching $2M$ however large the cluster grows (section 5.6.1 develops the comparison).
3. **Common pitfall**: A frequent misconception is that ring AllReduce makes every worker exchange data with every other worker. In ring AllReduce, each worker communicates with neighbors according to a schedule, and the per-node communication volume is approximately constant as the cluster grows. The scaling pressure comes from latency steps, topology, and the gradient volume per step, not from every worker opening a distinct stream to every other worker.

For standard synchronous data-parallel stochastic gradient descent (SGD), gradient synchronization supplies the shared gradient needed for one optimization trajectory. Replicas that apply different local gradients can diverge. To reproduce the intended global-batch update, they must also start from matching parameters and optimizer state, weight local contributions correctly, and apply the same update rule. The AllReduce primitive provides the reduction; convergence follows from the optimization assumptions, not from communication alone. Floating-point reduction order can still change low-order bits.

In general, a model with P parameters contributes $2P$ bytes per data-parallel replica when all gradients are communicated in BF16 (2 bytes per parameter); stored weight precision alone does not determine that choice. For the 70-billion-parameter model in the definition above, that complete logical gradient is 140 GB, and a ring moves nearly twice its input buffer size per rank.[2] If a replica spans multiple ranks through model parallelism, each corresponding-shard group reduces its own portion rather than this full buffer.

Because every bandwidth term in this chapter scales linearly with payload size, the worked derivations use this compact 70B gradient (or the chapter's reference gradient buffer first priced in the energy estimate of section 5.2) and vary rank count to isolate how each schedule scales. Section 5.7.1 then prices the full 175B-parameter running instance of Archetype A (GPT-4/Llama-3) on this book's reference cluster.

At large scale (hundreds of billions of parameters across thousands of GPUs), gradient synchronization can dominate training step time unless collective communication is systematically optimized. Understanding how to manage that overhead requires modeling data movement as physics rather than as an abstract API call.

5.2 The Physics of Data Movement

Whenever a gradient shard leaves an accelerator, three limits bound its trip before any collective algorithm runs. The speed of light sets a propagation floor[3] of roughly 5 µs per kilometer in fiber. High-frequency signal loss limits how far copper traces can reach before the link must convert to optics (the bandwidth-distance product of section 3.3.2). Host-kernel networking stacks add tens of microseconds of software copy overhead, unless kernel-bypass transports such as remote

[2] **Ring AllReduce**: The algorithm dates to the HPC collective-communication literature; Patarasuk and Yuan analyzed bandwidth-optimal AllReduce algorithms for workstation clusters, Baidu's Andrew Gibiansky popularized ring AllReduce for deep learning in February 2017, and Horovod plus NVIDIA NCCL helped make it a common distributed-training primitive (Patarasuk and Yuan 2009; Gibiansky 2017; Sergeev and Balso 2018; Jeaugey 2017).

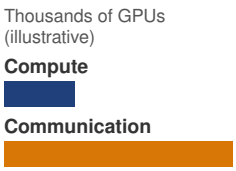

Exposed gradient synchronization can dominate step time at scale.

[3] **The speed-of-light constraint**: Light travels through optical fiber at approximately 200,000 km/s, or 200 meters per microsecond. For a rack-to-rack fiber path of 100 meters, propagation alone contributes approximately 500 ns in one direction. Faster switches can reduce other latency components, but reducing this propagation delay requires a shorter path or a different propagation medium.

direct memory access (RDMA) and GPUDirect RDMA strip per-message initiation down to a few microseconds.

These constraints govern different parts of a communication budget. Propagation latency sets a floor for every message regardless of size. Link bandwidth caps throughput once a transfer is large enough to fill the pipe. Protocol overhead adds a fixed per-message cost that penalizes fine-grained exchanges unless a kernel-bypass transport removes it. RDMA's zero-copy path[4] avoids staging each transfer through host memory and the kernel's socket buffers, keeping that per-message cost low. Because the physical fabric and its software stack fix these limits, the collective schedule determines how many times the job pays them.

Collective communication also pays an energy price for every bit moved across those links, which the schedule must respect as firmly as latency and bandwidth. Moving data consumes energy across storage arrays, interfaces, serializers, transceivers, and switches. A transfer that starts in local SRAM can stay on-package in high-bandwidth memory (HBM), cross to an intra-node peer GPU, or traverse intermediate switches to a remote node. Each step widens the physical accounting boundary that an energy estimate must include (figure 5.1). Each boundary activates distinct physical components: on-package links, intra-node crossbars, serializers, optical transceivers, and switch fabrics. Distance alone does not determine energy per bit, and an endpoint-only estimate cannot be compared directly with a complete fabric-path estimate. Locality remains useful because it can avoid transfers and active interfaces; the energy saving must be evaluated under the same workload, utilization, and accounting boundary.

↩ Kernel-bypass RDMA, which strips the per-message software overhead, and GPUDirect RDMA, which removes the host-memory bounce, are established in section 3.4.1.

4 **Zero-copy communication**: GPUDirect RDMA lets the NIC read gradients directly from accelerator memory, so a 140 GB gradient never passes through host DRAM, whereas a staged path writes and then rereads it there (280 GB of host-memory traffic per send).

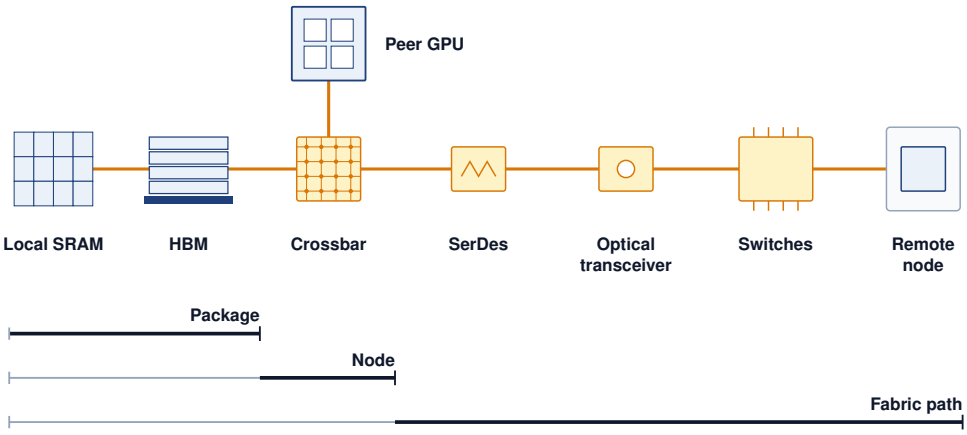

Figure 5.1: Data-Movement Accounting Boundaries: Successively broader paths involve different storage and communication components. The schematic identifies what an energy estimate must include; it assigns no universal distance law or fixed energy multiplier. Compare energy per transferred byte only after specifying endpoints, interfaces, switches, and utilization.

The chapter's reference gradient buffer of 1 GB stands for one large reduction, far smaller than a full model gradient but large enough that per-byte transfer, not per-message startup, governs its cost; later examples return to it by name. Reducing it around an illustrative 10,000-rank ring sends approximately twice the payload from every rank. At a fabric-path transfer cost of 35 pJ/bit, a representative midpoint of the range typical of switched fabrics rather than a measurement of any particular one, this gives about 5.6 kJ per reduction summed over all the ranks, paid again at every synchronous step. The usable rule is that communication energy is the bytes an algorithm sends multiplied by the energy per bit of every boundary those bytes cross. A byte that leaves the package or the node crosses boundaries an on-chip operand does not, which can raise its energy by orders of magnitude. An algorithm that moves fewer bytes, or a placement that keeps more of them inside the node, therefore saves energy as directly as it saves time.

The share of the step's total energy that communication represents must be measured at the same accounting boundary as the compute it accompanies. The fabric-path figure charges the serializers, transceivers, and switches between two endpoints. A whole-system estimate that also charges

endpoint memory reads, host software, and idle network power assigns each byte crossing the network a much larger cost, so two per-byte figures can be compared only after their boundaries match.

Across this section, propagation delay, link bandwidth, per-message software overhead, and energy per bit have bounded single transfers in isolation. Which collective operations a training step actually invokes, and how often they pay those limits, depends on how the workload partitions its parameters, activations, and tokens across devices.

↪ The energy cost of moving a byte across memory, storage, and network boundaries, and why movement can outweigh arithmetic, is developed in section 15.4.2.2.

5.3 Parallelism Strategies and Their Traffic

Once a training job splits its batch, layers, or parameters across devices, no rank holds all the state needed to finish a step on its own. A synchronous data-parallel worker computes gradients from its local batch slice, but its optimizer cannot update the weights until every replica's gradients have been combined. Some of those reductions can start as backward layers finish, while earlier layers are still computing. Tensor ownership determines where each tensor must go, and gradient or activation readiness determines when each transfer can start.

How the model is partitioned dictates which tensors cross device boundaries and which ranks must exchange them. For Archetype A (GPT-4/Llama-3), dense gradient contributions are combined within each data-parallel group; in a hybrid layout, that group contains the matching parameter shard from each replica. AllReduce serves the replicated-gradient baseline. For Archetype B (DLRM at Scale),[5] lookup requests travel to embedding-shard owners and the resulting vectors return to the workers that need them. Mixture-of-experts (MoE) routing similarly sends activations to the ranks hosting selected experts and returns their outputs. These destination-specific exchanges motivate AllToAll.

At large scale, these parallelism strategies can be combined. A 3D parallelism layout combines data parallelism, tensor parallelism, and pipeline parallelism, creating distinct communication groups whose exchanges may overlap in time. In a node-local tensor-parallel example, layer reductions use NVLink, pipeline stages exchange activations and their gradients point-to-point across nodes for each microbatch (a slice of the batch), and data-parallel groups reduce matching parameter shards over InfiniBand. Other placements change those physical paths while retaining the logical communication requirements.

Each primitive operates on a different process group, a subset of the total GPU population that participates in that particular collective. Classical Message Passing Interface (MPI) uses a communicator as the object that names such a participating set; its communicator size is the number of ranks in that set. MPI libraries already selected collective algorithms based on communicator size and message size (Thakur et al. 2005). GPU communication libraries inherit that algorithm-selection problem and add the further challenge of coordinating overlapping process groups without creating contention between concurrent collectives.

Each parallelism strategy fixes the data it must move, the collective primitive that moves it, and the term most likely to limit it (table 5.1), and with them which participants must exchange data. The collective primitives that follow make each exchange's contract precise, the pricing model then charges it per round, and the collective algorithms provide its execution schedule.

[5] **DLRM (Deep Learning Recommendation Model)**: Meta's 2019 architecture for click-through rate prediction uses sharded embedding tables and exchanges the resulting sparse lookups across workers, an AllToAll whose limiting term section 5.8 develops.

↩ The three lighthouse archetypes and the C^3 regimes each one stresses are established in section 1.6.1.

Table 5.1: Parallelism-to-Collective Mapping: How each parallelism strategy becomes a traffic pattern. Pipeline parallelism alone relies on point-to-point transfers rather than a collective.

Parallelism strategy	What must move	Primary primitive	Primary constraint
Data Parallelism	Full gradient, averaged across replicas	AllReduce	Bandwidth (large payloads)
FSDP/ZeRO-3	Parameter shards in, gradient shards out	AllGather + ReduceScatter	Bandwidth, high frequency
Tensor Parallelism	Partial layer activations within the node	AllReduce, AllGather	Critical path (NVLink bandwidth and latency)
Sequence Parallelism	Activations sharded along the sequence	AllGather + ReduceScatter	Bandwidth (activation exchange)
Pipeline Parallelism	Activations to the next stage	Point-to-Point (Send/Recv)	Latency (microbatch dependencies)

Parallelism strategy	What must move	Primary primitive	Primary constraint
Expert Parallelism (MoE)	Tokens to the ranks hosting their experts	AllToAll	Startup, injection/bisection, or skew
Embedding Sharding (DLRM)	Embedding lookups to and from shard owners	AllToAll	Granularity, fabric capacity, or skew

In this table, expert parallelism refers to the mixture-of-experts architecture pattern, in which each token activates only a few experts. Data parallelism and fully sharded data parallel (FSDP), which divides each replica's parameters, gradients, and optimizer state across the data-parallel ranks, generate large reductions and gathers that reward byte-efficient rings and hierarchical schedules. Tensor parallelism, by contrast, generates smaller, more frequent messages directly on the critical path, so low software overhead and fast node-local links dominate its cost. Pipeline parallelism also sits on the critical path, because each stage waits for the previous stage's microbatch, but its stages often exchange activations across nodes, so the size of those activations sets its transfer regime. Expert parallelism and sharded embedding tables produce destination-specific traffic whose bottleneck shifts among message granularity, endpoint injection and reception, fabric bisection, oversubscription, routing, and skew.

The table's rows describe data center parallelism, so they leave out the federated Archetype C. Across all three lighthouse archetypes, the dominant exchange and its bottleneck differ.

Lighthouse 5.1: Communication archetype patterns

A gradient's communication pattern depends on the system's objective function and constraint regime. Each lighthouse archetype faces a distinct communication challenge, and the techniques developed in this chapter map to those challenges differently, as table 5.2 summarizes.

Table 5.2: Archetype-to-Collective Mapping: Each lighthouse archetype produces a characteristic traffic pattern, but the limiting term depends on message shape, placement, and the physical fabric.

Archetype	Primary collective	Candidate limiter	Optimization strategy
Archetype A (GPT-4/Llama-3)	AllReduce	Bandwidth (β)	Hierarchical AllReduce; rail-optimized placement
Archetype B (DLRM at Scale)	AllToAll	Granularity, fabric capacity, or skew	Message coalescing; topology-aware routing; shard balancing
Archetype C (Federated MobileNet)	Client uploads to an aggregator (asynchronous)	Uplink bandwidth and availability	Aggressive quantization; error feedback

The table separates traffic pattern from limiting resource. A dense reduction, a destination-specific exchange, and an intermittent device update require different interventions even though all three serve a training workload. Hierarchical AllReduce, AllToAll schedule selection, and compression with error feedback each receive their own treatment later in the chapter.

A traffic map names which pattern a workload generates, but it does not specify what each rank must hold when the call completes. Before any of those interventions can be priced, each exchange needs a precise contract.

5.4 Collective Operation Primitives

Sending a gradient tensor to a peer does not specify which workers must contribute or where the final result belongs. Collective operations supply that group-level contract, using a vocabulary standardized by MPI and inherited by modern ML communication libraries (Message Passing Interface Forum 2015). The four contracts sketched at the start of the chapter differ in what each rank holds afterward (figure 5.2). AllReduce leaves an identical sum (Σ) on every rank. AllGather

leaves every rank with the concatenation of all shards. ReduceScatter sums across ranks but leaves each rank one reduced shard (Σ_i). AllToAll performs a distributed matrix transpose, routing each rank's j-th chunk to rank j. Two further primitives complete the set and serve as building blocks for these four: Broadcast, in which rank 0 sends its data to all ranks, and Reduce, in which all ranks aggregate into rank 0.

↪ The bandwidth and latency cost terms of AllReduce, AllGather, ReduceScatter, AllToAll, and Broadcast, under ring, tree, and pairwise-exchange algorithms, are developed in section D.2.

Definition 5.2: Collective operation

Collective operation is a distributed communication pattern in which every rank in a process group invokes a matching operation to aggregate, broadcast, or redistribute data. The operation's contract specifies the output at each rank; it does not require simultaneous entry and does not, by itself, impose a general-purpose barrier before or after the call (Message Passing Interface Forum 2015).

1. **Significance**: Because the contract names only each rank's result, the schedule that carries it decides how communication demand spreads as the group grows. A schedule that funnels every contribution through one endpoint makes the communication term $T_{\text{comm}}(N)$ of the Fleet Law (equation 1.3) grow with N, while a balanced schedule keeps each rank's bytes nearly constant; section 5.6.1 quantifies the gap.
2. **Distinction**: Unlike point-to-point communication, a collective defines one operation over a process group. Ranks may enter at different times. In Broadcast or Reduce, which center on one rank, some algorithms allow a rank to complete before another has entered, whereas an AllReduce result needs every rank's contribution. Correctness depends on matching the operation sequence and argument contract across ranks, not on barrier semantics.
3. **Common pitfall**: Treating a non-blocking collective as complete when its initiation call returns. The returned request or stream event represents work in flight. A dependent consumer must wait for completion, and a late rank can extend that dependency even though the collective is not a barrier.

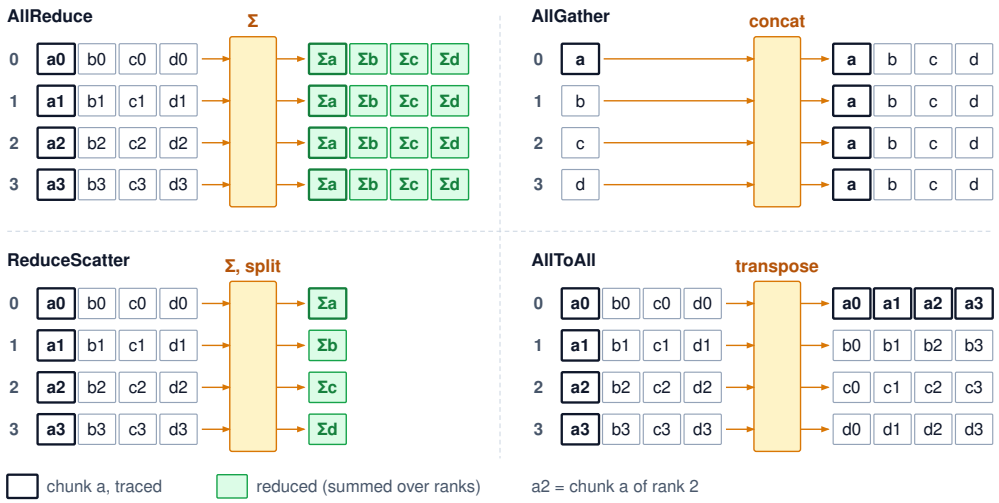

Figure 5.2: Four Central Collective Primitives: Standardized patterns for group communication in distributed systems. In each panel the four rows are ranks 0 to 3, read left to right from each rank's input, through the collective, to what each rank holds afterward. (1) AllReduce: all aggregate and all receive the result. (2) AllGather: all send to all and concatenate. (3) ReduceScatter: all aggregate and results are distributed in shards. (4) AllToAll: every process sends unique data to every other process. The two foundational primitives Broadcast (rank 0 sends to all) and Reduce (all aggregate to rank 0) compose the patterns shown here and are defined in the surrounding prose.

The primitive choice is the first scaling decision because it fixes which ranks must coordinate and which result each must obtain. Once that contract is chosen, the algorithm and placement determine how the resulting traffic crosses the fabric.

5.4.1 The six core primitives

Calling a full AllReduce when each rank's optimizer updates only a $1/N$ parameter shard moves twice as many bytes as the step needs, because half the work goes into assembling full tensors that every rank immediately discards. Choosing the right primitive therefore means matching the contract in table 5.3 to the exact slice the next computation consumes.

The table's rows run from the root-centered patterns (Broadcast and Reduce), through the three in which every rank contributes to and receives part or all of one shared tensor (AllReduce, AllGather, and ReduceScatter), to AllToAll, in which every pair of ranks exchanges distinct data. That order describes routing rather than cost. ReduceScatter and AllGather each move half of AllReduce's bytes because each is one of the two phases into which AllReduce decomposes.

The cost column previews the algorithms derived in section 5.6, using N for the participating rank count, M_{coll} for the full tensor, and M for the per-rank AllToAll input payload. For AllReduce, the full tensor has the same byte size as each rank's input payload, the convention used in the following algorithm models. Two schedules appear there: a ring passes chunks between neighbors, whereas a simple full-message tree reduces up a binary tree and broadcasts back down it, forwarding the whole tensor at each of its $2\log_2 N$ steps. For Broadcast and Reduce, the table counts startup rounds and bytes per tree edge separately. Each traverses one tree depth, so the elapsed transfer cost accumulates across its levels. The Tree AllReduce derivation combines those two traversals.

Table 5.3: The Six Collective Primitives: What each primitive computes, its primary use, and its communication cost, ordered from root-centered patterns (Broadcast, Reduce) to the fully destination-specific AllToAll. The order describes routing, not cost.

Primitive	What it does	Primary use case	Communication cost
Broadcast	One sender transmits data to all receivers.	Distribute initial model weights from rank 0 at startup, across all parallelism strategies.	$\mathcal{O}(\log N)$ startup rounds; $\mathcal{O}(M_{\text{coll}})$ bytes per tree edge. Elapsed byte-transfer time also grows with tree depth.
Reduce	Aggregate data from all workers (sum, min, max) to a single root.	Aggregate validation metrics such as loss and accuracy to a logging process.	$\mathcal{O}(\log N)$ startup rounds; $\mathcal{O}(M_{\text{coll}})$ bytes per tree edge. Elapsed byte-transfer time also grows with tree depth.
AllReduce	Aggregate from all workers, then distribute the result to all ($y_i = \sum_{j=0}^{N-1} x_j$ for all i).	Data parallelism: synchronize gradients so every GPU computes the same update.	Ring is bandwidth-optimal at $2\frac{N-1}{N}\frac{M_{\text{coll}}}{\beta}$ but $\mathcal{O}(N)$ latency; the simple full-message Tree model gives $\mathcal{O}(\log N)$ latency with more byte-transfer work.
AllGather	Each worker's data goes to all; the result concatenates every input ($x_i \to [x_0, \ldots, x_{N-1}]$).	Sharded data parallelism (FSDP/ZeRO): collect sharded parameters before a forward or backward pass.	Ring $\frac{N-1}{N}\frac{M_{\text{coll}}}{\beta}$; resident data grows from an M_{coll}/N shard to the full M_{coll}.
ReduceScatter	Reduce across workers, but scatter so each keeps a distinct chunk (worker i receives the i-th block of $\sum x_j$).	Sharded data parallelism: reduce gradients while keeping them sharded to save memory.	Ring $\frac{N-1}{N}\frac{M_{\text{coll}}}{\beta}$; the same ring byte-transfer term as AllGather, with reduction.
AllToAll	The most general pattern: worker i sends a distinct chunk to each worker j, a distributed matrix transpose.	MoE routing tokens to experts; DLRM exchanging embedding lookups across workers.	Per-worker payload $\frac{N-1}{N}M$; the limiter can arise from startup, injection/reception, bisection, oversubscription, routing, granularity, or skew.

Ring achieves AllReduce's bandwidth-optimal cost while the simple full-message Tree model trades bandwidth for logarithmic latency (Patarasuk and Yuan 2009; Thakur et al. 2005). AllReduce itself also decomposes into ReduceScatter followed by AllGather. The decomposition is operational as well as mathematical, since Ring AllReduce's first phase is a ReduceScatter and its second phase is an AllGather. ReduceScatter performs arithmetic, so it does not undo an AllGather even though it

shrinks buffers where AllGather expands them. Full Zero Redundancy Optimizer (ZeRO) sharding and PyTorch FSDP use the pair the same way, gathering parameter shards with AllGather before a module needs them and reducing gradients with ReduceScatter after backward so each rank retains only its shard for the optimizer step (Rajbhandari et al. 2020). Sharding parameters, gradients, and optimizer state across workers reduces persistent model-state memory roughly by $N\times$, while temporary full-parameter buffers and prefetching determine peak memory. That memory saving is developed in section 6.3.3; the communication it adds is priced in section 5.4.2.

The same AllGather/ReduceScatter pair also appears in sequence parallelism, which splits activations along the sequence dimension. There an AllGather reconstructs the activation view needed for a local operation, and a later ReduceScatter redistributes the result along that dimension. A primitive's semantics stay fixed even when the tensor dimension being sharded changes, so every schedule developed for the pair applies to both uses.

Beyond defining data-movement dependencies, collective operations carry explicit interface and numerical contracts. At the interface level, every rank in a process group must issue collectives in a compatible sequence; mismatched operation orders can deadlock or pair unrelated buffers. Operations also enforce shape and type contracts. AllReduce requires matching element counts and datatypes across all ranks, whereas variable-size exchanges require explicit per-destination counts and displacements. These invariants form the contract between the parallel decomposition and the communication runtime. The model decides which tensor slices each rank owns, and the collective contract guarantees that those slices can be recombined deterministically.

Reduction order introduces a second, numerical contract. Floating-point addition is not associative, so ring, tree, and switch-assisted reductions (in which network switches do the adding, developed in section 5.7.3) may accumulate the same values in different orders and produce different low-order bits (Goldberg 1991). A correct AllReduce therefore guarantees the specified reduction over all contributions, not bitwise identity across algorithms, rank counts, or hardware. Reproducibility-sensitive workloads must fix the algorithm and process-group configuration or use a reproducible accumulation scheme; ordinary training should treat small order-dependent differences as numerical behavior while still checking for larger rank disagreement.

AllToAll scales differently from AllReduce with cluster size, and section 5.8 develops the bound that governs it. With the six contracts defined, every row of table 5.1 now names a primitive whose output and traffic pattern are precise. The sharded data-parallel row repays closer study, because sharding changes how often those primitives run rather than which ones run.

5.4.2 FSDP communication patterns

Sharding every parameter, gradient, and optimizer-state tensor across N data-parallel ranks saves memory, but it replaces one end-of-step reduction with a stream of collective calls inside the forward and backward passes. Under fully sharded data parallelism (Rajbhandari et al. 2020), each rank holds only $1/N$ of a layer's weights. It must call AllGather before the layer's forward pass to assemble the full weights. After the backward pass, it calls ReduceScatter to leave each rank the gradient for its own shard, and then frees the temporary full weights.

Trading resident memory for on-demand assembly multiplies the collective call count. With one communication group per layer, a schedule that holds each layer's gathered weights until its backward pass finishes issues $2N_L$ collective calls per step (one forward AllGather and one backward ReduceScatter per layer) where data parallelism can reduce its whole gradient in one AllReduce. To save peak activation memory, full-resharding FSDP instead frees the full weights right after each layer's forward pass and runs a second AllGather before that layer's backward pass, raising the count to about $3N_L$ calls per step. The byte volume changes far less. Because each AllGather and ReduceScatter over a large group moves nearly M_{coll} bytes per rank across all layers combined, where M_{coll} is the model-sized payload without optimizer state, the $2N_L$-call schedule sends roughly the AllReduce's $2M_{\text{coll}}$ per rank and the resharded $3N_L$-call schedule sends about $3M_{\text{coll}}$ ($1.5\times$) when parameters and gradients travel at the same width. Different parameter and reduction datatypes shift that ratio.

The call count carries its own price, because every call pays a startup cost. An effective startup cost α_{eff} contributes approximately $2N_L\alpha_{\text{eff}}$ per step before overlap, or $3N_L\alpha_{\text{eff}}$ with resharding. FSDP is therefore more sensitive to α_{eff} than a single whole-gradient reduction even though its

bytes barely change. The counts describe one grouping policy, since FSDP units need not coincide with layers and data parallelism can also split its reduction into several calls to enable overlap. FSDP implementations mitigate the startup cost through prefetching and communication-stream pipelining (issuing collectives on a stream separate from computation), but the relevant α_{eff} must come from a matched collective benchmark rather than a wire-latency specification.

↪ The ZeRO stages and the reshard-after-forward FSDP schedule, with the memory each saves and the communication each adds, are developed in section 6.3.3.

Selecting the right collective algorithm at each invocation requires distinguishing transfer cost from dependency exposure. An FSDP AllGather sits on the critical path, because a layer cannot compute until its full parameters arrive. Its transfer can nevertheless be bandwidth-bound when the parameter tensor is large. Prefetching can hide part of that gather, while a ReduceScatter can overlap with the remaining backward computation. Neither dependency alone selects Ring or Tree; message size, group size, topology, and the available overlap window jointly determine the useful schedule.

The primitives specify what must be communicated, and sharding shows how often a training step invokes them. Choosing a schedule for each invocation first requires a price for every round of an exchange.

5.5 Pricing Collectives by Rounds and Chunks

A gradient bucket, the group of gradient tensors reduced in one call, can reach tens of megabytes, large enough to saturate a single network link. Yet when a thousand-GPU ring divides that bucket into a thousand shards, each hop moves only tens of kilobytes and can spend more time starting the transfer than sending its bytes. Pricing a collective therefore requires separating the fixed time to initiate each message from the rate at which bytes move once in transit.

The α-β model $T(n) = \alpha + n/\beta$ prices a transfer over a single link (principle 5), separating a fixed startup cost (α) from a per-byte transfer cost (n/β) (Hockney 1994). A collective pays that price at every step of its schedule, so the two terms, multiplied by the schedule's round count and traffic factor (the multiple of the tensor each rank sends), determine algorithm selection across the fleet.

↩ The α-β model for a single link and its latency-dominated and bandwidth-dominated regimes are established in section 3.4.4.

For collectives, the parameters must describe the effective path. Latency (α) is the startup cost paid at every step of a schedule, so it includes recurring software overhead, PCIe traversal, and network switching time; one-time communicator initialization is a separate setup cost. Because the ranks in a step send concurrently, elapsed time counts one startup per step rather than one per rank. Bandwidth (β) is the sustained per-direction transfer rate. The critical message size $n^* = \alpha \cdot \beta$, where the two terms are equal, carries over unchanged. It must, however, be compared against each step's chunk (M/N in an N-rank ring) rather than the whole tensor M. That is why a large bucket can still fall below the inter-node crossovers in table 5.4. Below n^*, a step is latency-bound; above it, bandwidth-bound.

↪ Worked message-size regimes and crossover calculations for this model are developed in section D.1.

The crossover explains two recurring optimizations. For InfiniBand NDR at $\alpha \approx 1.5$ µs and $\beta \approx 50$ GB/s, the knee sits at $n^* \approx 75$ KB, the value section 3.4.4 computed for a single transfer. A 4 KB token-routing message is far below n^*, so the startup cost dominates. If 100 ready messages have compatible destinations and can share one transfer, fusing them into a 400 KB message reduces the modeled cost from approximately $100\alpha = 150$ µs to $\alpha + n/\beta \approx 9.5$ µs. That is roughly a 16× improvement before packing costs. More generally, N_{msg} serialized messages of size n/N_{msg} can cost up to $N_{\text{msg}}\times$ as much as one message of size n when $n/N_{\text{msg}} \ll n^*$. Fusion is the first optimization, and it motivates framework gradient buckets and destination-aware token coalescing.

Compression is the second, and the crossover tells where it pays. A 140 GB gradient tensor sits far above n^*, so shrinking its bytes shrinks the dominant n/β term directly, provided encoding and aggregation costs are counted. Even keeping only the largest 1 percent of values (Top-k sparsification, developed in section 5.9) leaves 1.4 GB of values before their indices are added, still far above n^*, so the saving remains a byte saving. Compressing a 4 KB token-routing message, by contrast, offers little, because its unchanged α term already dominates.

Across the typical data center interconnects tabulated in table 5.4, intra-node NVLink stays latency-bound up to several hundred kilobytes, whereas the inter-node fabrics cross over below roughly 100 KB. A message between those crossover sizes can therefore be latency-bound inside a node but bandwidth-bound once it travels between nodes.

Table 5.4: Interconnect Performance Parameters: Effective per-message launch latency (α, including the software stack), per-direction link bandwidth (β, the one-way deliverable rate that enters $T(n) = \alpha + n/\beta$; NVLink's datasheet 900 GB/s is the bidirectional total, so its β is the 450 GB/s one-way rate), and the computed crossover size $n^* = \alpha \cdot \beta$ at each row's α midpoint (for InfiniBand NDR, at the representative 1.5 µs used throughout this chapter), where communication transitions from latency-bound to bandwidth-bound.

Interconnect	Latency (α)	Bandwidth (β)	Critical size (n^*)
NVLink 4.0 (intra-node)	1–2 µs	450 GB/s	~0.7 MB
InfiniBand NDR 400 Gbps	1–3 µs	50 GB/s (per port)	~75 KB
InfiniBand HDR 200 Gbps	2–5 µs	25 GB/s	~87.5 KB
PCIe Gen5 (GPU↔CPU)	2–5 µs	64 GB/s	~224 KB
Ethernet 100 Gbps (RoCE)	5–10 µs	12.5 GB/s	~93.7 KB

The NVLink and NDR rows show the roughly ninefold per-direction cliff of section 2.5.1, which section 3.1 traces across a decade of interconnect generations. Even on the lower-crossover inter-node fabric, a bucket far above n^* can still travel in chunks below it because a ring divides the bucket among its N ranks. Pricing one bucket at two group sizes shows when that happens and what it costs.

Napkin Math 5.1: The critical message size

Problem: A data-parallel job reduces a 25 MB gradient bucket with Ring AllReduce over InfiniBand NDR, with $\alpha = 1.5$ µs and $\beta = 50$ GB/s. The bucket is hundreds of times larger than the single-transfer knee. Does each step of the ring run in the bandwidth-bound regime on 64 GPUs, and on the 1,024-GPU reference cluster?

Math:

1. **Knee**: $n^* = \alpha \cdot \beta = 1.5 \times 10^{-6}$ s $\times\ 5 \times 10^{10}$ B/s $= 75$ KB.
2. **Chunk per step**: A ring sends M/N per step, 390.6 KB on 64 ranks (above the knee) but only 24.4 KB on 1,024 ranks (below it). The ring takes one fewer step than it has ranks in each of its two phases.
3. **Ring time** $2(N-1)\left(\alpha + \frac{M}{N\beta}\right)$: about 1.2 ms on 64 ranks, of which the byte term is 1.0 ms; about 4.1 ms on 1,024 ranks, where the byte term is still only 1.0 ms.

Systems insight: The regime belongs to the chunk, not to the tensor. The byte term barely moves as the group grows, but at 1,024 ranks each step carries a chunk below n^*, so startup dominates and the same bucket takes several times longer. Large language model (LLM) gradients are bandwidth-bound as tensors, yet a flat ring over a large group can still carry them in latency-bound steps. The levers then shift to larger buckets, fewer rounds, and hierarchy, which is why the algorithms of section 5.6 compete on round count as well as bytes.

For collectives, the two regimes map onto different levers. In the latency-bound regime ($n < n^*$), where fine-grained per-peer token-routing chunks, scalar reductions, and the small chunks of a large-group schedule sit, the levers are fusion, fewer schedule rounds, fewer hops, and a leaner software stack. In the bandwidth-bound regime ($n \gg n^*$), where LLM gradients and optimizer states sit, they are compression, a byte-efficient algorithm, and link aggregation across multiple network interface cards (NICs). A token-routing exchange can fall on either side, because its regime depends on the chunk sent to each peer, which shrinks as the group grows if each rank's payload stays fixed. Classifying each exchange, and each chunk within it, into one regime is the diagnostic skill the rest of the chapter applies.

Checkpoint 5.1: Per-chunk pricing diagnostics

Verify your understanding of per-chunk pricing regimes:

- ☐ **Chunk regime**: classify each ring step of a 25 MB gradient bucket reduced over 1,024 ranks, on links with $\alpha = 5$ µs and $\beta = 50$ GB/s, as latency-bound or bandwidth-bound.
- ☐ **Software overhead**: determine which type of parallelism benefits more when software overhead (α) is reduced by half, data parallelism or tensor parallelism.
- ☐ **Chunks under an upgrade**: explain why upgrading links from 100 Gbps to 400 Gbps at constant software latency can move a large ring's per-step chunks from the bandwidth-bound to the latency-bound side of $n^* = \alpha \cdot \beta$.
- ☐ **NIC scaling**: assess the True or False claim that, in the bandwidth-bound regime, doubling the number of NICs per node effectively doubles the transmission speed for large tensors.

The pricing model charges a schedule per round and per chunk, but it does not choose the schedule. AllReduce, the most performance-critical primitive, admits several schedules whose round counts and byte volumes differ sharply, and for a large gradient the choice can decide whether gradient synchronization can be hidden at all.

5.6 AllReduce Algorithms and Topologies

A correct AllReduce must give each of 1,000 GPUs the reduced sum of all 1,000 gradient contributions, subject to the floating-point reduction contract, and the algorithm that delivers it sets the cluster's scaling ceiling. The naive answer, in which every GPU sends to one central server whose single network link must carry every worker's gradient, fails first, which makes it the right place to start.

5.6.1 Naive approaches vs. the bandwidth bottleneck

Consider a naive implementation using a parameter server (a star topology), with a dedicated server in addition to the workers. All N workers send their gradients to that server; it sums them and sends the result back. The server must receive $N \times M$ bytes and send $N \times M$ bytes, so its bandwidth budget grows as $T_{\text{comm}} \propto N \times M/\beta$. At 1,000 GPUs, concentrating those transfers at one server creates a scaling bottleneck even if the arithmetic is fast.

This naive approach captures the pressure that early distributed ML frameworks had to manage. Parameter-server systems receive, aggregate, and redistribute model updates through a server tier, and practical implementations shard the parameter space to spread that load (Dean et al. 2012; Li et al. 2014). Sharding reduces the per-server traffic from $N \times M$ to $N \times M/K_{\text{srv}}$ (where K_{srv} is the number of servers). The server tier still grows as workers and model state grow, and it introduces additional complexity in parameter partitioning and consistency management.

The bottleneck comes from concentration, not from total traffic. Every AllReduce algorithm moves more total data as workers are added, because every worker's contribution must be combined and every worker must receive the result. A star, however, routes all of that traffic through one endpoint, so the server's time grows as $N \times M/\beta$. A ring spreads the same work so that each rank sends about $2M$ and its time stays near $2M/\beta$ however large N becomes. The ratio of the two is about $N/2$, so at 1,024 workers the server's byte time is roughly 500× the ring's, before startup, arithmetic, or contention are counted. Bandwidth optimality makes that comparison precise by measuring an algorithm's required byte movement against a lower bound under a stated communication model, which the next derivation supplies for AllReduce.

War Story 5.1: When the parameter server became the bottleneck (2017)

Context: In 2017, Alexander Sergeev and Mike Del Balso's team at Uber open-sourced Horovod as part of Uber's Michelangelo ML platform, where TensorFlow training jobs needed to scale across many GPUs without forcing every model author to rewrite large portions of training code (Sergeev and Balso 2018).

Mechanism: TensorFlow's stock distributed training relied on a parameter-server pattern that concentrated gradient traffic through a central tier. As worker count N grew, the central server link saturated, degrading scaling efficiency.

Impact: Scaling training beyond a single node resulted in severe throughput bottlenecks and required complex, model-specific parameter partitioning.

Fix: Uber adopted Baidu's ring AllReduce algorithm (Gibiansky 2017), wrapping bandwidth-optimal collective reduction behind a lightweight API that replaced parameter servers with ring communication.

Systems lesson: A collective algorithm is also a developer interface. Scaling improves when the communication primitive is both bandwidth efficient and easy enough that teams actually use it correctly.

5.6.2 The bandwidth-optimal lower bound

Before choosing a schedule to replace the parameter server, we need to know the minimum byte volume any point-to-point AllReduce must move when N ranks each start with an M-byte gradient and neither compression nor in-network switch reduction is used. Reducing one independent data item across N ranks and distributing its result requires at least $2(N-1)$ item transfers in aggregate (Patarasuk and Yuan 2009). The count has two halves. Combining N contributions that start on N different ranks takes at least $N-1$ transfers, because each transfer combines at most two partial sums and so reduces the number of separate partial sums by at most one. Delivering the finished value to the other $N-1$ ranks then takes at least one transfer each. For an M-byte tensor, the minimum aggregate sent volume is therefore $2(N-1)M$ bytes. At least one rank must send the average share, $2(N-1)M/N$ bytes; a balanced algorithm assigns that share to every rank.

Take β as each rank's available send bandwidth, with equal receive bandwidth that can operate simultaneously. Dividing the required sent volume by β gives the full-duplex bandwidth lower bound:

$$T_{\text{bandwidth}}^{\min} = \frac{2(N-1)}{N} \cdot \frac{M}{\beta}$$

As N grows large, the bandwidth lower bound approaches $2M/\beta$, independent of N (principle 11). Ring AllReduce matches the required byte volume under this model; startup and local reduction still add time (Patarasuk and Yuan 2009; Thakur et al. 2005). On a 64-GPU cluster synchronizing the chapter's 1 GB reference buffer over 50 GB/s links, the ring's byte term is about 39.4 ms, establishing the physical floor against which every concrete AllReduce schedule that follows must be judged.

5.6.3 Ring AllReduce

To hit that $2(N-1)M/N$ byte floor without funneling traffic through a central server, Ring AllReduce arranges nodes in a logical loop ($0 \to 1 \to \ldots \to N-1 \to 0$) where each node sends to its successor while receiving from its predecessor at the same time. Splitting the tensor into equal chunks spreads the required byte volume across every rank and reaches the bandwidth lower bound in the ideal endpoint model (Patarasuk and Yuan 2009).

The algorithm splits the vector of size M into N chunks and then alternates communication and local reduction in two phases, a ReduceScatter phase (often called scatter-reduce) followed by an AllGather phase. Algorithm 5.1 states the mechanics before the figure and trace make the same dataflow concrete.

Algorithm 5.1 Ring AllReduce

Require: N ranks in a logical ring; each rank i starts with tensor $\mathbf{g}_i$ of M bytes, split into N chunks
Ensure: every rank holds $\mathbf{g} = \sum_{i=0}^{N-1} \mathbf{g}_i$

1: assign the ring order $0 \to 1 \to \ldots \to N{-}1 \to 0$
2: **for** $N-1$ ReduceScatter steps **do**
3: each rank sends one chunk (or partial sum) right, receives one from the left, and adds it into the matching local sum
4: **end for**
5: each rank now owns one fully reduced chunk of $\mathbf{g}$
6: **for** $N-1$ AllGather steps **do**
7: each rank forwards a reduced chunk around the ring and stores the chunks it receives
8: **end for**
9: **return** the complete reduced tensor $\mathbf{g}$, held on every rank

Across the $2(N-1)$ send/receive rounds each rank moves M/N bytes per round, for $2(N-1)M/N$ bytes sent and the same received. The bandwidth term therefore reaches the information-theoretic lower bound, while the latency term grows as $2(N-1)\alpha$. On a four-GPU ring with wrap-around links, the state progresses in two phases (figure 5.3). In row (a), $N-1=3$ ReduceScatter steps circulate gradient chunks A through D until each GPU holds exactly one fully reduced partition (green ΣA through ΣD). In row (b), $N-1=3$ subsequent AllGather steps forward those finished shards around the ring until every device holds the complete sum.

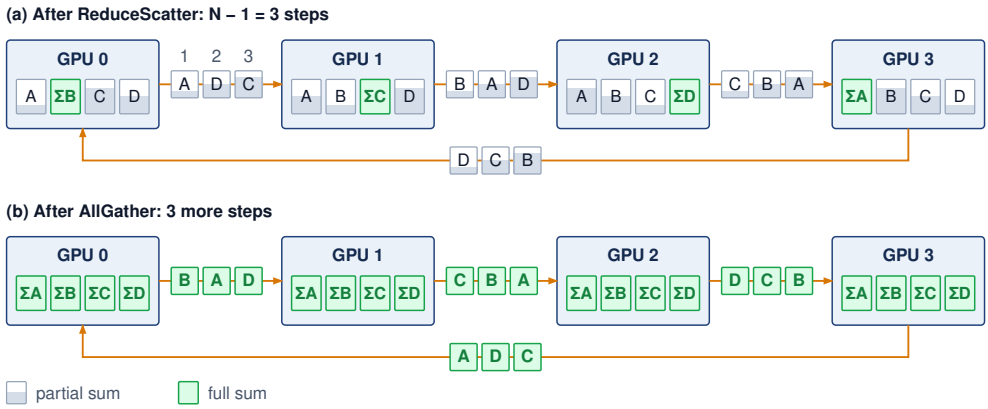

Figure 5.3: Ring AllReduce Data Flow: Four GPUs in a logical ring, each gradient split into chunks A to D; the wrap-around link closes the ring from GPU 3 back to GPU 0. (a) After the $N-1$ ReduceScatter steps, in which every GPU sends one chunk to its neighbor and adds the chunk it receives, each GPU holds one fully reduced chunk. (b) After $N-1$ AllGather steps that forward those finished chunks around the same ring, every GPU holds the full sum. Every link carries data in every step, so each GPU sends $2\frac{N-1}{N}M$ bytes in total, while the $2(N-1)$ steps each pay one startup latency α.

Logical ring edges may share physical links, and startup, local reduction, or contention can leave transfer capacity unused, so byte optimality does not guarantee peak physical-link utilization. To see how the logical schedule itself reaches that byte lower bound, consider 4 GPUs reducing a vector of 4 elements by summation. Each GPU i starts with a local gradient vector $\mathbf{g}_i = [a_i, b_i, c_i, d_i]$ split into 4 chunks (one per GPU), and the algorithm proceeds through two phases.

Example 5.1: Ring AllReduce: Two-phase example (4 GPUs)

Scenario: Four GPUs ($0 \to 1 \to 2 \to 3 \to 0$) perform a Ring AllReduce on local vectors $[1,5,3,7]$, $[2,6,4,8]$, $[3,7,5,9]$, and $[4,8,6,10]$ to compute the element-wise sum $[10,26,18,34]$.

Diagnosis: The trace has to show three properties of the schedule, namely that it takes $2(N-1)=6$ steps, that every GPU sends and receives in every step, and that each step moves only one chunk. Ring AllReduce splits each vector into $N=4$ chunks a, b, c, d. In ReduceScatter step s ($s=1,2,3$), GPU i sends its running copy of chunk $(i-s+1)$ mod 4 to GPU $(i+1)$ mod 4, which adds it to its own copy. After the third step, each GPU owns one fully reduced chunk. In AllGather step s, GPU i forwards the finished chunk it completed or received most recently to GPU $(i+1)$ mod 4, which stores it without arithmetic. In each of the six steps, every GPU receives one chunk and updates the value it holds (table 5.5).

Table 5.5: Ring AllReduce Trace: The chunk each of four GPUs receives in each of the six steps. In the three ReduceScatter steps the GPU adds the incoming chunk to its running partial sum; in the three AllGather steps it stores an already finished sum.

Step	GPU 0	GPU 1	GPU 2	GPU 3
ReduceScatter 1	d: 17	a: 3	b: 13	c: 11
ReduceScatter 2	c: 14	d: 25	a: 6	b: 21
ReduceScatter 3	b: 26	c: 18	d: 34	a: 10
AllGather 1	a: 10	b: 26	c: 18	d: 34
AllGather 2	d: 34	a: 10	b: 26	c: 18
AllGather 3	c: 18	d: 34	a: 10	b: 26

Every cell in every row is filled, so every GPU receives (and, symmetrically, sends) one chunk in every step, and after the sixth step every GPU holds [10, 26, 18, 34]. The steps are also chained, because the chunk a GPU sends next is the one it has just received. A GPU with a slow link therefore holds up every step, and in bandwidth-bound steps the whole collective slows by about the same factor as that link.

Systems lesson: Ring AllReduce achieves the optimal byte volume under the stated model: each GPU sends $2\frac{N-1}{N}M$ total bytes, approaching $2M$ as N grows, while the startup-round count continues to grow. Physical-link utilization also depends on placement and runtime overhead.

Checkpoint 5.2: Ring AllReduce mechanics

Verify your understanding of bandwidth-optimal reduction:

- ☐ **Step count**: calculate the number of sequential steps required for a ring of N GPUs to complete the full AllReduce.
- ☐ **Bandwidth optimality**: explain why Ring AllReduce is bandwidth-optimal, including how many times each byte of the total message M traverses the network per GPU.
- ☐ **Straggler impact**: determine the impact on total collective completion time when one GPU in the ring is 50 percent slower than the others.
- ☐ **Send/receive duality**: assess the True or False claim that every GPU is both a sender and a receiver in every step of Ring AllReduce.

Generalizing the six-step trace of example 5.1 to N ranks yields the closed-form cost model for those checkpoint questions. Each node sends and receives $\frac{M}{N}$ bytes in each of the $2(N-1)$ steps.

$$T_{\text{ring}} = \underbrace{2(N-1)\alpha}_{\text{Latency Term}} + \underbrace{2\frac{N-1}{N}\frac{M}{\beta}}_{\text{Bandwidth Term}}$$

The bandwidth term approaches $2M/\beta$ as $N \to \infty$, reaching the full-duplex lower bound. Each rank sends nearly $2M$ bytes across the two phases and receives the same amount concurrently. In contrast, the latency term scales linearly with N. Across 10,000 nodes, 19,998 sequential startup rounds accumulate substantial latency, which explains why Ring AllReduce becomes uncompetitive for small payloads. For a multi-gigabyte gradient buffer at moderate rank count, however, the bandwidth term dominates completely, as the 70B calculation in notebook 5.2 confirms.

Napkin Math 5.2: AllReduce cost for a 70B model

Problem: Consider a communication-only baseline for a 70B parameter model across 64 GPUs connected by InfiniBand NDR (50 GB/s per port). The bandwidth term barely depends on rank count, so this smaller group reproduces the bandwidth-only estimate from the opening definition while keeping the latency term visible. Each rank contributes a complete logical gradient buffer communicated in BF16 (2 bytes per parameter). This calculation isolates collective cost without establishing that the training state fits on each GPU. How long does one AllReduce take?

Step 1: Size the gradient payload. Each rank in this baseline contributes the complete logical gradient tensor of $7 \times 10^{10} \times 2$ bytes $= 140$ GB. A hybrid layout that partitions the model would instead reduce corresponding gradient shards within smaller groups and requires its own calculation.

Step 2: Apply the Ring AllReduce bandwidth formula.

$$T_{\text{bandwidth}} = 2 \cdot \frac{N-1}{N} \cdot \frac{M}{\beta}$$

Substituting: $T_{\text{bandwidth}}$ = 1.96875 × 140 GB / 50 GB/s ≈ 5,512.5 ms.

Step 3: Add the latency term.

$$T_{\text{latency}} = 2(N-1) \cdot \alpha$$

Substituting: $T_{\text{latency}} = 126 \times 1.5$ µs = 0.2 ms.

Step 4: Total communication time. Total: $T_{\text{ring}} \approx$ 5,512.5 ms + 0.2 ms ≈ 5,512.7 ms.

Systems insight: The modeled full-buffer AllReduce takes over five seconds under the stated link and rank assumptions, and almost all of it is the bandwidth term. Its effect on step time depends on how much remains exposed relative to useful computation. When it is fully exposed, halving it saves about 2.8 s per step; across a 100,000-step campaign on these 64 GPUs, that is roughly 4,900 GPU-hours of reserved accelerator time. Large-scale training systems can launch reductions for output-side gradients while input-side layers are still computing. Partitioned training also changes each rank's payload and participating group, so this complete-gradient budget must be recomputed for the actual layout.

5.6.4 Tree AllReduce

Ring AllReduce's byte efficiency comes with many startup rounds. As the group grows, the $2(N-1)\alpha$ term can exceed the byte-transfer term for a fixed payload; the crossover depends on the measured path parameters. Tree AllReduce reduces the number of dependent rounds. In the simple reduce phase, pairs combine contributions and forward partial sums toward a root, taking $\log_2 N$ steps for a balanced power-of-two group. The broadcast phase distributes the result in another $\log_2 N$ steps. This shorter dependency chain motivates the comparison before bandwidth costs are counted.

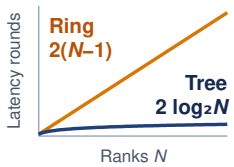

Ring latency grows with N; tree stays logarithmic.

The simple full-message tree loses parallel activity as contributions converge. A ring activates all N logical edges in each round. In a pairwise tree reduction, the first round has $N/2$ senders and $N/2$ receivers, so every rank is still active. The next round has $N/4$ senders and $N/4$ receivers, and the final reduction exchange involves two ranks while the remaining $(N-2)$ have completed their upstream contribution. Sending the full message through successive levels adds byte-transfer

time along the dependency path. This is a limitation of the simple schedule, not every tree-based implementation or every physical link mapping.

The resulting time complexity depends on the specific tree variant. A simple reduce-then-broadcast tree has logarithmic latency but nonuniform traffic and root/interior bottlenecks. Recursive doubling, in which partners at doubling distances exchange the entire tensor in each of $\log_2 N$ rounds, also has a logarithmic round count and incurs $\mathcal{O}(\log N)$ full-message bandwidth per rank. The following simplified model, which uses the simple tree's counts, captures the latency advantage and the potential bandwidth penalty of tree-like collectives:

$$T_{\text{tree-like}} \approx \underbrace{2\log_2 N \cdot \alpha}_{\text{Latency}} + \underbrace{2\log_2 N \frac{M}{\beta}}_{\text{Bandwidth penalty in full-message exchange variants}}$$

The latency term is logarithmic ($\mathcal{O}(\log N)$). For 1,024 nodes, Ring needs 2,046 steps while Tree needs only 20. That roughly 100× reduction in steps makes Tree attractive for small, frequent collectives across large process groups, such as scalar reductions, control messages, or small-batch activation reductions. The bandwidth term carries the penalty, because tree-like algorithms can underutilize links or concentrate traffic on interior links, and recursive-doubling-style full-message exchanges add a $\log_2 N$ bandwidth factor. These penalties make naive tree variants a poor choice for data parallelism's large gradient AllReduce, where bandwidth efficiency determines whether the network is saturated or partially idle.

For large-scale data parallelism, this bandwidth penalty decides the comparison whenever the implementation repeatedly exchanges full messages or creates root/interior bottlenecks. On the 64-GPU cluster of section 5.6.2, the simple tree forwards the whole tensor once per level on the way up and once per level on the way down, $2\log_2 64 = 12$ full-message transfers. Its byte time is about 240 ms against the ring's 39.4 ms, or 6.1× the ring's byte time. The difference of roughly 200 ms per step would accumulate to about 5.6 hours over a 100,000-step run if fully exposed, although actual wall-clock savings also depend on startup, placement, and overlap. For a 70B-parameter model's 140 GB gradient AllReduce across 1,000 GPUs, a bandwidth-inefficient tree-like variant can erase the latency advantage that motivated it. Optimized double-tree algorithms exist precisely to keep logarithmic latency without imposing that naive bandwidth penalty.

Small activation reductions provide a useful case for fewer startup rounds. For an eight-GPU group communicating a 1 MB tensor, the simple model reduces Ring's $2(8-1)\alpha = 14\alpha$ to Tree's $2\log_2(8)\alpha = 6\alpha$, but Tree's byte term remains about 3.4× Ring's. Whether the saved startups outweigh those extra byte transfers depends on the measured latency and bandwidth. Tensor-parallel collectives also have a scheduling constraint independent of this comparison: a layer that needs the reduced activation must wait for it. A critical-path dependency does not by itself imply a latency-dominated transfer or select a tree algorithm.

5.6.5 Recursive halving-doubling (butterfly)

The simple full-message Tree reduces startup rounds but repeats more byte-transfer work than Ring. Recursive halving-doubling, sometimes called the Butterfly algorithm, combines logarithmic rounds with the optimal byte volume under the flat-network model. It pairs ranks exactly as recursive doubling does, but where recursive doubling exchanges the whole tensor in every round, halving-doubling halves the exchanged region in each ReduceScatter round and doubles it in each AllGather round. That difference removes the $\log_2 N$ byte penalty. For a power-of-two group, it uses $2\log_2 N$ rounds: $\log_2 N$ ReduceScatter rounds followed by $\log_2 N$ AllGather rounds. A partner-bit index k selects the XOR pairing with mask 2^k. The ReduceScatter trace below visits those bits in increasing order as messages shrink; the matching AllGather reverses that order as the assembled result grows.

In the ReduceScatter phase (first $\log_2 N$ rounds), GPU i partners with GPU $i \oplus 2^k$ (XOR of indices). Because the two partners made identical retention choices in all earlier rounds, both enter round k holding partial sums for the exact same $M/2^k$-byte region of the tensor. They split that active region in half, retain complementary halves of size $M/2^{k+1}$, and send the other half to each other to add into the half each retained. Across the $\log_2 N$ rounds, the bytes each GPU sends form a geometric series $M/2 + M/4 + \cdots + M/N = \frac{N-1}{N}M$, leaving each GPU with a distinct M/N-byte slice of the

fully reduced result. AllGather traverses the partner bits in reverse, exchanging completed regions of doubling size $M/N, 2M/N, \ldots, M/2$ until every rank holds the full result.

Summing the $2\log_2 N$ startups and the two symmetric geometric byte series gives the recursive halving-doubling cost:

$$T_{\text{butterfly}} = 2\log_2 N \cdot \alpha + 2\frac{N-1}{N} \cdot \frac{M}{\beta}$$

The schedule combines logarithmic latency ($\mathcal{O}(\log N)$, like Tree) with bandwidth-optimal data movement ($2(N-1)/N \cdot M/\beta$, like Ring). The catch is that it requires nonneighbor communication (GPU i must communicate with GPU $i \oplus 2^k$, which may be physically distant), and it requires N to be a power of two. For clusters where N is not a power of two, additional complexity is needed to handle the irregular cases.

Recursive halving-doubling is used in MPI-style collective implementations and is useful as a theoretical point in the latency-bandwidth trade-off (Thakur et al. 2005). NVIDIA's GPU collective backend, the NVIDIA Collective Communications Library (NCCL), offers implementation families that include ring, tree-derived, hierarchical, NVLink/NVSwitch-aware, and pattern-aware choices, with exact names and availability changing by version and topology (Jeaugey 2017; NVIDIA 2026c). Whatever a particular release calls them, practical communication libraries select algorithms through a topology-aware cost model because nonlocal communication patterns can create contention on shared network links that erodes theoretical advantages.

The partner distance is where topology enters. Consider the ReduceScatter phase for 8 GPUs placed four per node on two nodes, GPUs 0–3 on the first and GPUs 4–7 on the second, which proceeds in $\log_2(8) = 3$ rounds. In round $k = 0$, each GPU i partners with GPU $i \oplus 1$, pairing neighbors: (0,1), (2,3), (4,5), (6,7). They exchange half their data and reduce, and the lower-numbered GPU of each pair keeps the first half of the tensor while its partner keeps the second half. In round $k = 1$, the distance doubles: GPU i partners with GPU $i \oplus 2$, creating pairs (0,2), (1,3), (4,6), (5,7). GPUs 0 and 2 both hold the first half, so they split it between them, and GPUs 1 and 3 likewise split the second half. In round $k = 2$, the distance doubles again: GPU i partners with GPU $i \oplus 4$, creating pairs (0,4), (1,5), (2,6), (3,7).

Physical placement determines the cost of these logical pairings (figure 5.4). Rounds $k = 0$ and $k = 1$ pair GPUs within a node, but in round $k = 2$ every GPU on the first node exchanges with one on the second, so that round crosses the node boundary. On eight-GPU nodes the same crossing would appear one round later, at $k = 3$ in a 16-rank group. Keeping consecutive ranks on one node confines the large exchanges to fast local links and leaves only the smallest ones on the slower fabric. In a separate 1,024-rank example, the final ReduceScatter round pairs GPU i with GPU $i \oplus 512$. If the two rank halves occupy different racks, those exchanges share the inter-rack fabric. The algorithm's logarithmic rounds and optimal byte volume therefore do not establish its runtime advantage; placement, shared paths, and the changing chunk size in each round must enter the comparison.

5.6.6 Double binary tree

Where recursive halving-doubling requires power-of-two groups and long-distance XOR pairings across fabric cuts, a tree schedule can follow physical topology more flexibly, provided it overcomes the single-tree bandwidth bottleneck. In a single binary tree, half the ranks sit at the leaves and only send (during reduction) or only receive (during broadcast), leaving half of each leaf's bidirectional link capacity idle while interior nodes do double duty. NCCL resolves this imbalance with a double binary tree (Jeaugey 2017; NVIDIA 2026c), constructing two independent binary trees with complementary interior-node roles and running both simultaneously so each carries half the payload.

The complementary-tree construction balances endpoint work by making ranks that act as interior nodes in one tree act as leaves in the other. With eight ranks, build the first tree by labeling the nodes of a binary tree in the order of an in-order traversal (left subtree, then node, then right subtree). In a balanced seven-node tree labeled from 0, in-order traversal puts the leaves at the even ranks 0, 2, 4, and 6 and the interior nodes at 1, 3, and 5. In one simple construction, the eighth rank, 7, becomes the root, with rank 3, the root of the seven-node tree, as its only child. A single child is enough, because rank 7 joins the interior only so that every rank is interior in exactly one of the two trees. The odd ranks 1, 3, 5, and 7 thus form the first tree's interior and the even ranks are its leaves.

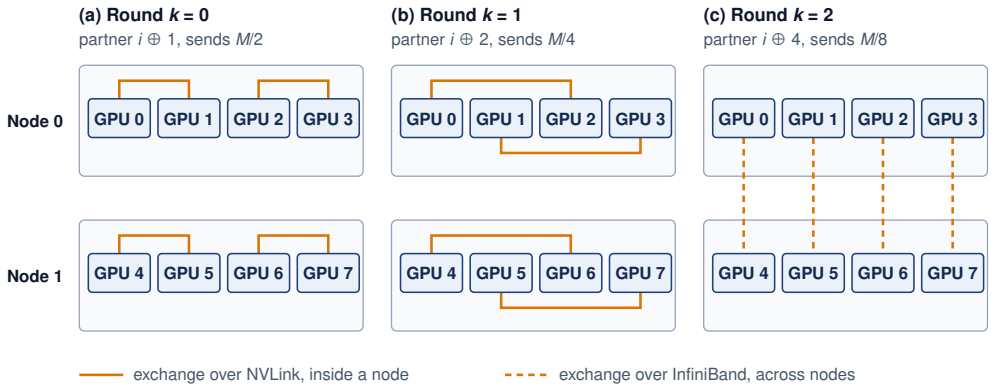

Figure 5.4: Halving-Doubling Partners across a Node Boundary: The ReduceScatter phase of recursive halving-doubling on 8 GPUs placed four per node. In round k, GPU i exchanges with GPU $i \oplus 2^k$ and sends half of the region it still holds, so the exchanged volume halves each round. (a) and (b) The first two rounds pair GPUs inside a node and use NVLink. (c) The last round pairs every GPU with one on the other node, so the round with the smallest message is the only one that crosses the slower fabric. Which round crosses depends on how many ranks each node holds.

The second tree shifts every rank by one position, so the even ranks become its interior and the odd ranks its leaves (figure 5.5). Each tree carries half the tensor, $M/2$, and the two halves never meet.

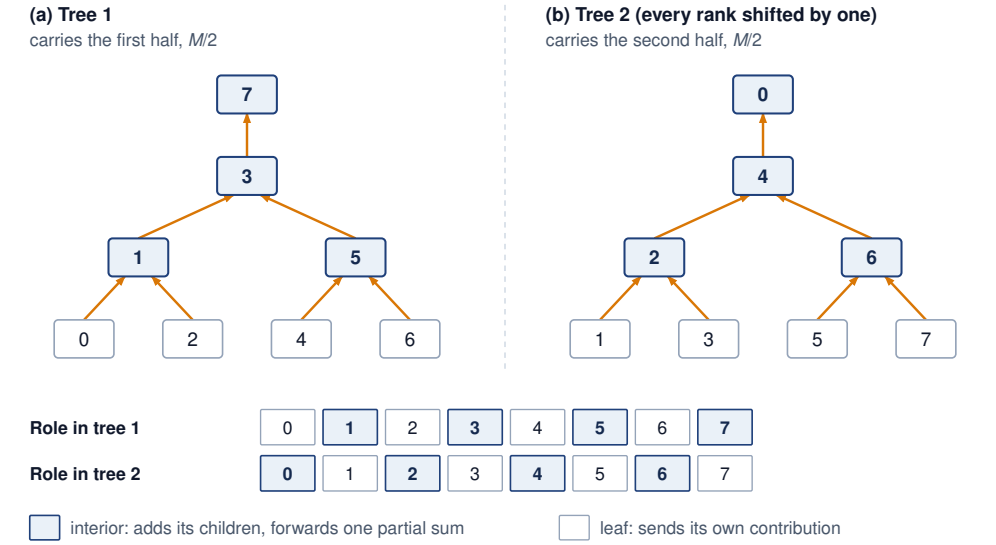

Figure 5.5: Complementary Roles in a Double Binary Tree: Two reduction trees over eight ranks, each carrying half of the tensor toward its root. (a) The first tree places the odd ranks at its interior and the even ranks at its leaves. (b) The second tree shifts every rank by one, which swaps the roles. The strip below shows the result: every rank is an interior node in exactly one tree and a leaf in the other, so it forwards one combined half and sends one half of its own, about M in the reduce phase. The broadcast phase runs the same trees in the opposite direction.

Two further steps bring the byte term down from the simple tree's $2\log_2 N \cdot M/\beta$ to about $2M/\beta$. First, each tree pipelines its half in small chunks, so while one chunk moves up from the leaves, the next is already leaving them and every level works at once. Tree depth then adds startup rounds and a short pipeline fill, not repeated full-message transfers. Second, count what one rank sends during the reduce phase. As an interior node of one tree it combines the inputs of its (at most two) children and forwards one $M/2$ stream to its parent. As a leaf of the other tree it sends its own $M/2$ contribution. It therefore sends about M in total. The broadcast phase mirrors that traffic, for about

$2M$ per rank across both phases, the same per-rank volume as the ring. This logical balance does not guarantee that every physical link stays busy or that both trees avoid shared fabric bottlenecks; rank placement and routing still matter.

The combined performance is approximately:

$$T_{\text{double-tree}} \approx 2\log_2 N \cdot \alpha + \frac{2M}{\beta}$$

The ideal bandwidth term approaches $2M/\beta$ while the startup-round count remains logarithmic. This combination makes the double binary tree a useful candidate across message sizes that penalize either the ring's many rounds or the simple tree's repeated full-message transfers. Complementary endpoint roles do not ensure disjoint physical paths, so topology-aware tree construction seeks a good placement, and measurements must establish the achieved rate under contention.

5.6.7 The algorithm crossover point

Plotting the three cost models for a 256-GPU group on a 200 Gbps fabric, over message sizes from 1 KB to 10 GB, shows the simple Tree and Ring curves intersecting near 1 MB (figure 5.6). At small sizes, fewer startup rounds favor the tree schedules; at large sizes, the simple full-message Tree pays more byte-transfer time. The Double Binary Tree curve includes illustrative overhead factors, so it approaches a comparable bandwidth-dominated regime without representing a measured library result. The candidate schedules differ in their latency and bandwidth terms, in whether they reach the bandwidth lower bound, and in the constraint each one places on the group (table 5.6).

Table 5.6: AllReduce Algorithm Comparison: Each algorithm occupies a different point in the latency-bandwidth trade-off space. Ring and Butterfly achieve bandwidth optimality with different latency characteristics, while Double Binary Tree combines logarithmic depth with a near-optimal byte term. The Tree row refers to the simple full-message model, not every tree implementation.

Algorithm	Latency	Bandwidth	Bandwidth optimal?	Constraint
Ring	$\mathcal{O}(N)\alpha$	$2\frac{N-1}{N}\frac{M}{\beta}$	Yes	None
Simple Tree	$\mathcal{O}(\log N)\alpha$	$\mathcal{O}(\log N)\frac{M}{\beta}$	No	None
Butterfly	$\mathcal{O}(\log N)\alpha$	$2\frac{N-1}{N}\frac{M}{\beta}$	Yes	$N = 2^m$
Double Tree	$\mathcal{O}(\log N)\alpha$	$\approx \frac{2M}{\beta}$	Near-optimal	Complementary tree construction

The intersection near 1 MB in figure 5.6 marks where the modeled Ring and simple Tree have equal total cost. It is an algorithm crossover, distinct from the startup-versus-transfer crossover for one message. Setting $T_{\text{ring}} = T_{\text{tree}}$ and solving for payload M identifies the point where Ring's lower byte cost compensates for its additional startup rounds:

$$T_{\text{ring}} = 2(N-1)\alpha + 2\frac{N-1}{N}\frac{M}{\beta}$$

$$T_{\text{tree}} = 2\log_2 N \cdot \alpha + 2\log_2 N \cdot \frac{M}{\beta}$$

For large N, $(N-1) \approx N$ and $\frac{N-1}{N} \approx 1$, so the ring pays a latency term proportional to N while moving each byte close to the bandwidth limit:

$$T_{\text{ring}} \approx 2N\alpha + \frac{2M}{\beta}, \quad T_{\text{tree}} \approx 2\log_2 N \cdot \alpha + \frac{2\log_2 N \cdot M}{\beta}$$

The tree has the opposite shape: logarithmic startup cost, but each message byte moves through more stages. Setting the estimates equal isolates the message size where the cheaper latency path stops compensating for the extra bandwidth work:

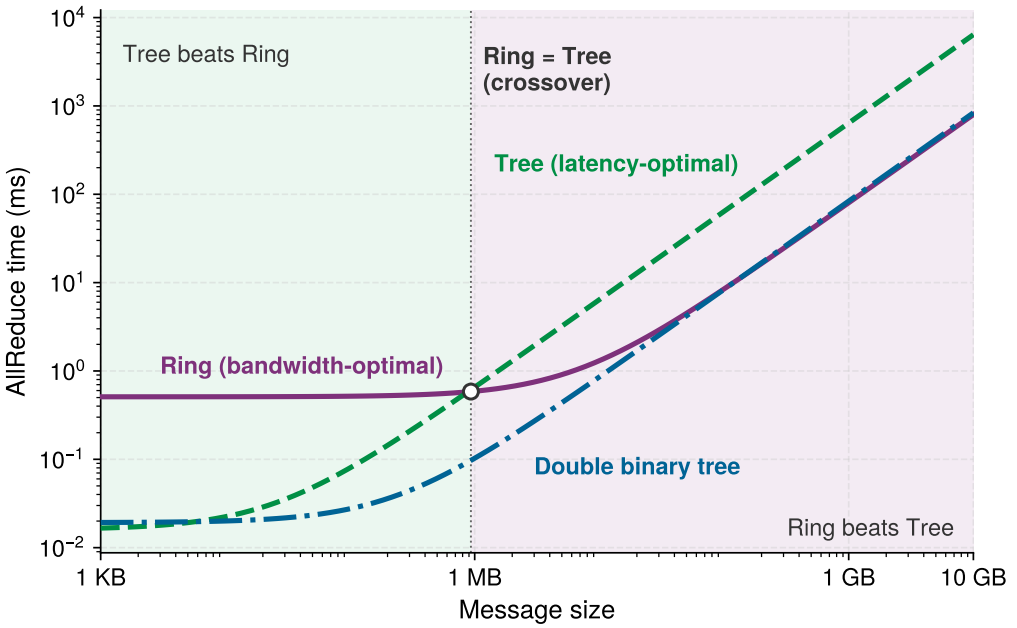

Figure 5.6: Algorithm Crossover in Collective Communication: Below about 1 MB the simple tree's fewer startup rounds make it faster than the ring; above that size the ring's smaller byte volume wins, and the double binary tree tracks the faster of the two across most sizes. The curves are modeled AllReduce times for a 256-GPU group on a 200 Gbps interconnect with $\alpha = 1$ µs, over messages from 1 KB to 10 GB. The double binary tree curve applies modest latency and bandwidth overhead factors (1.2× and 1.05×) to its ideal cost, so it illustrates the model rather than a measured library result.

$$2N\alpha + \frac{2M}{\beta} = 2\log_2 N \cdot \alpha + \frac{2\log_2 N \cdot M}{\beta}$$

$$\frac{2M}{\beta} - \frac{2\log_2 N \cdot M}{\beta} = 2\log_2 N \cdot \alpha - 2N\alpha$$

$$\frac{2M}{\beta}(1 - \log_2 N) = 2\alpha(\log_2 N - N)$$

Since $N \gg \log_2 N$ for large clusters, $(\log_2 N - N) \approx -N$ and $(1 - \log_2 N) \approx -\log_2 N$, yielding:

$$\frac{2M}{\beta}(-\log_2 N) \approx 2\alpha(-N)$$

$$\boxed{M_{\text{crossover}} \approx \frac{N \cdot \alpha \cdot \beta}{\log_2 N}}$$

The log-aware crossover formula is the large-N approximation to the exact crossover $M_{\text{crossover}}(N)$ derived in section D.3.1. Within this two-algorithm model, Ring wins above the crossover because its byte savings exceed its extra startup cost; the simple Tree wins below it because the startup saving is larger. The crossover rises with N because each added rank adds Ring startup rounds faster than it adds Tree byte transfers. It rises with α because each of those rounds costs more, and with β because each byte saved is worth less time. A bucket that favors Ring on a small group can therefore favor Tree on a larger one. Optimized trees and topology-dependent rates require separate comparisons.

On the reference cluster, with the chapter's effective NDR $\alpha = 1.5$ µs and $\beta = 50$ GB/s across all 1,024 GPUs, the estimate gives $M_{\text{crossover}} \approx \frac{1{,}024 \times 1.5\ \mu\text{s} \times 50\ \text{GB/s}}{\log_2 1{,}024} = \frac{1{,}024 \times 1.5\ \mu\text{s} \times 50\ \text{GB/s}}{10} \approx 7.7$ MB. A gradient bucket of tens of megabytes or the full model gradient therefore sits on the ring side of the crossover, while a reduction of a few hundred kilobytes sits on the tree side. The section's two other

configurations change the group and the fabric, which moves the crossover. A smaller 256-GPU group on a slower 200 Gbps fabric with $\alpha = 1$ µs moves the curves' intersection down to about 1 MB (figure 5.6). A 64-GPU group on a still slower, higher-latency fabric places a 1 MB buffer almost exactly on the crossover (notebook 5.3).

This crossover separates only Ring from the simple Tree. Butterfly pays Ring's byte term with logarithmic startup rounds (table 5.6), so for a power-of-two group on a flat network the model gives Ring no win over Butterfly. Ring's case rests on what that model leaves out. Ring runs on any group size, and its exchanges stay between neighbors, whereas Butterfly's long-distance pairings can contend on shared links (section 5.6.5).

NCCL's algorithm selection is more nuanced than a two-way split. The double binary tree offers logarithmic latency with near-optimal bandwidth, so it competes with both Ring and Tree across a broad range of message sizes. The library also weighs the number of channels (parallel communication streams), the network topology, and whether a hop is inter-node or intra-node. The effective selection logic resembles a multi-way decision indexed by message size, GPU count, and topology type rather than a single crossover point. The crossover formula nevertheless explains the shape of those choices, because every candidate trades startup rounds against bytes moved. The model therefore names a candidate worth benchmarking rather than proving the library wrong, and an override is retained only after measured collective duration and exposed training-step time confirm it.

Training frameworks decide which payloads reach that analysis. PyTorch DistributedDataParallel (DDP) groups gradient tensors into buckets (section 5.10.3), while FSDP groups parameter gathers and gradient reductions according to its sharding units. A bucket size alone does not locate a crossover, because rank count, effective startup cost, and available bandwidth also enter the comparison. NCCL's algorithm selection must therefore be evaluated for the actual invocation and topology rather than inferred from a nominal bucket-size range, so the worked comparison that follows states its group size, payload, and fabric.

Napkin Math 5.3: The ring vs. tree crossover

Problem: A cluster synchronizes a 1 MB buffer across 64 GPUs. The network has latency $\alpha =$ 10 µs and bandwidth $\beta = 10$ GB/s. Which algorithm performs better: Ring or Tree?

Math:

Latency:

1. **Ring Latency**: $2(N-1)\alpha = 2 \times 63 \times 10$ µs = 1260 µs.
2. **Tree Latency**: $2(\log_2 N)\alpha = 2 \times 6 \times 10$ µs = 120 µs.

Bandwidth (note the difference):

3. **Ring Bandwidth**: $2\frac{N-1}{N}\frac{M}{\beta} \approx 2 \times \frac{1\text{ MB}}{10\text{ GB/s}}$ = 196.9 µs (optimal: each rank sends about $2M$ in total, the lower bound).
4. **Tree Bandwidth**: $2\log_2 N \cdot \frac{M}{\beta} = 12 \times \frac{1\text{ MB}}{10\text{ GB/s}}$ = 1200 µs (each level sends full message).

Summing the latency and bandwidth terms gives each algorithm's total time (table 5.7).

Table 5.7: Ring vs. Tree AllReduce Timing for a 1 MB Buffer on 64 GPUs: Latency, bandwidth, and total-time terms with $\alpha = 10$ µs and $\beta = 10$ GB/s.

Algorithm	Latency	Bandwidth	Total
Ring	1260 µs	196.9 µs	1456.9 µs
Tree	120 µs	1200 µs	1320 µs

Tree wins, but only by 9 percent. For this 1 MB message, the payload is near the crossover point.

Systems insight: Ring's latency term (10.5× Tree's) nearly balances Tree's bandwidth term (6.1× Ring's). The log-aware crossover estimate predicts $M_{\text{crossover}} \approx N\alpha\beta/\log_2 N = 64 \times 10\ \mu\text{s} \times 10\ \text{GB/s}/6 \approx 1.1$ MB. At 1 MB, the payload is just below crossover, so Tree wins. At 10 MB, Ring would dominate.

The crossover point sits inside the range that gradient reductions can traverse as bucket sizes change. In this example, the 1 MB AllReduce lies near crossover, whereas a 10 MB DDP bucket is already above it. A fused bucket of 50–300 MB or the full 140 GB BF16 gradient of a 70B parameters model favors Ring over the simple Tree model even more strongly. These conclusions apply to the stated AllReduce configuration. An MoE token-routing payload of the same size still requires an AllToAll schedule, and optimized tree or hierarchical implementations require their own cost comparison.

Checkpoint 5.3: AllReduce algorithm selection

Verify your understanding of Ring vs. Tree AllReduce trade-offs:

- ☐ **Ring vs. Tree**: explain why Ring AllReduce achieves bandwidth-optimal communication (approaching $2M/\beta$) while the simple full-message Tree model incurs $\mathcal{O}(\log N)$ bandwidth overhead.
- ☐ **Crossover point**: calculate $M_{\text{crossover}}$ for a 256-GPU cluster with $\alpha = 5$ µs and $\beta = 50$ GB/s, then determine whether a 10 MB gradient buffer should use Ring or Tree.
- ☐ **ReduceScatter trace**: trace through the ReduceScatter phase of Ring AllReduce for 3 GPUs with a 3-element vector and verify that each GPU ends with the correct partial sum.
- ☐ **Latency limit**: explain why Ring AllReduce's $\mathcal{O}(N)$ latency hurts small, frequent collectives across a large process group, and why that penalty is modest inside an eight-GPU tensor-parallel group.

The schedules in this section differ in the startup rounds they pay and the bytes they send, and a crossover estimate turns that difference into a choice for a given message size and group size. That choice holds only as far as the model's assumptions about the network hold.

5.7 Hierarchical Communication

The AllReduce cost models in section 5.6 assume a flat network with uniform endpoint bandwidth, whereas real clusters expose the tiered bandwidth hierarchy of section 2.5.1. Because intra-node links are abundant and fast while inter-node links are scarce and expensive, hierarchical communication shrinks payloads locally before they cross the scarce tier, decomposing AllReduce into a local ReduceScatter, a cross-node reduction, and a local AllGather. It also aligns traffic with the physical fabric, whether switches aggregate packets in the network, rail-optimized trees separate GPU lanes, or direct-connect pods wire accelerators into multidimensional tori.

5.7.1 Hierarchical AllReduce

On a DGX H100 node, each GPU's local NVLink capacity is 450 GB/s per direction (the 900 GB/s datasheet figure is the bidirectional total), while its one inter-node port provides 50 GB/s over InfiniBand. These are endpoint capacity budgets, not a promise of uniform achieved rate to every peer. Intra-node connectivity and simultaneous traffic still matter. That hierarchical network motivates a schedule that assigns work to each tier according to its usable capacity.

A flat ring advances at the rate of its slowest edge, and every ordering of a multi-node ring crosses node boundaries, so it can leave local NVLink capacity unused while loading scarce inter-node links. Hierarchical AllReduce explicitly partitions the global operation into three phases that align communication volume with the fabric tiers (figure 5.7).

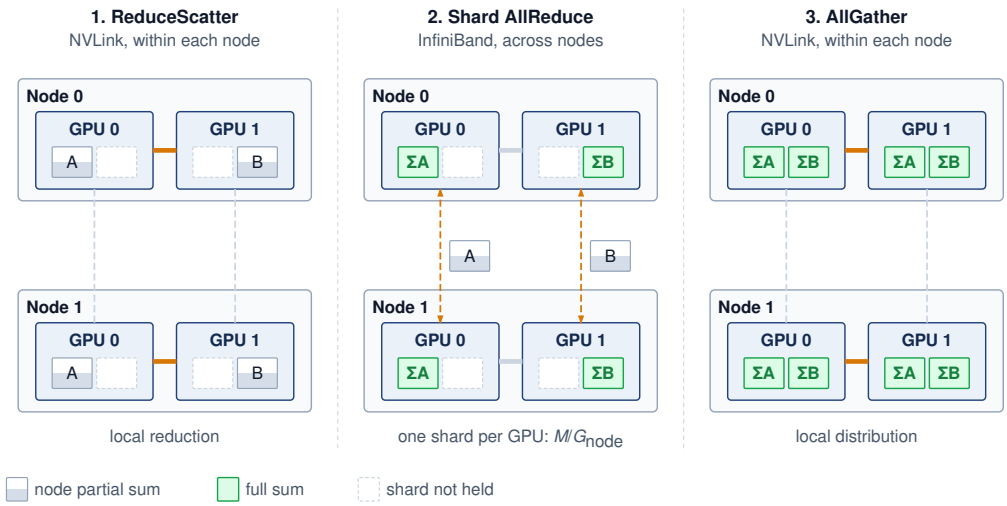

Figure 5.7: Hierarchical AllReduce: Three-phase decomposition that exploits the intra-node/inter-node bandwidth gap. Phase 1 performs intra-node aggregation over NVLink (an NVLink ring acting as a ReduceScatter that leaves each GPU holding a node-shard partial sum), leaving a payload of M/G_{node} on each of the G_{node} GPUs per node. Phase 2 runs an AllReduce across corresponding GPUs using InfiniBand. Phase 3 completes intra-node distribution over NVLink (the AllGather that returns each GPU to the full result). The schematic shows two representative ranks per node; the same decomposition applies to G_{node} ranks.

With G_{node} GPUs in each of N_{nodes} nodes, the three phases form a single causal pipeline. First, an intra-node ReduceScatter over NVLink performs a local reduction among the G_{node} GPUs in each node, leaving each GPU with a $1/G_{\text{node}}$ shard of partially reduced data while consuming only abundant intra-node bandwidth. Next, each corresponding-rank group applies an inter-node AllReduce over InfiniBand to its M/G_{node}-byte input shard, where the chosen collective algorithm dictates the bytes transmitted across the network fabric. Finally, an intra-node AllGather over NVLink redistributes the globally reduced shards within each node, completing the collective without consuming additional inter-node bandwidth.

Only Phase 2 (Shard AllReduce) loads the slower inter-node InfiniBand links (figure 5.7). In this middle phase, each corresponding-rank group reduces an M/G_{node}-byte shard across N_{nodes} nodes. A ring in that group sends $2(N_{\text{nodes}}-1)M/(N_{\text{nodes}}G_{\text{node}})$ bytes per participating GPU and receives the same amount. Summing across the node's G_{node} GPUs gives $2(N_{\text{nodes}}-1)M/N_{\text{nodes}}$ bytes sent per node. Smaller per-port payloads improve elapsed time when the node has independent ports and sufficient fabric capacity; sharding does not create bandwidth on a shared uplink. An 8-node budget quantifies the bandwidth multiplier this decomposition provides when those ports operate in parallel.

Napkin Math 5.4: The hierarchical bandwidth multiplier

Problem: A cluster has 8 nodes, each with 8 GPUs (64 total). The system must AllReduce the chapter's 1 GB reference buffer. Assume each GPU has its own full-rate inter-node port, one rail per GPU as on a DGX H100 (rails are the subject of section 5.7.4.1), with no oversubscription. The flat baseline is a single ring over all 64 ranks. How much faster is Hierarchical AllReduce than that flat ring?

Flat ring AllReduce (ignoring hierarchy):

- Each GPU sends ~2 GB total (the bandwidth-optimal Ring AllReduce formula).
- Every edge of a ring carries the same traffic in every step, so the whole ring advances at the rate of its slowest edge. Every ordering puts at least one fabric-crossing edge on each node, even the node-contiguous ordering that crosses the fabric only once per node.
- **Effective Bandwidth**: Limited by the slowest link = 50 GB/s (InfiniBand).
- **Time**: $\approx 2 \times 1$ GB$/50$ GB/s = 40 ms (bandwidth term dominates).

Hierarchical AllReduce (3-phase decomposition):

1. **Intra-node ReduceScatter**: Each GPU sends 875 MB at 450 GB/s per direction $\rightarrow$ ~1.94 ms
2. **Inter-Node AllReduce**: Each GPU AllReduces a 1 GB$/8$ = 125 MB shard over InfiniBand; Ring moves roughly $2(N_{\text{nodes}} - 1)/N_{\text{nodes}}$ times that payload, and after adding the small latency term the phase takes ~4.40 ms at 50 GB/s
   (The eight shard groups run concurrently, each on its own GPU's port, so the node's eight ports together carry the full payload in parallel.)
3. **Intra-node AllGather**: Each GPU receives 875 MB at 450 GB/s per direction $\rightarrow$ ~1.94 ms
4. **Total time**: $\approx$ 1.94 ms + 4.40 ms + 1.94 ms = 8.3 ms

Systems insight: Hierarchical AllReduce achieves a 4.8× speedup by reducing the inter-node payload from 1 GB to 125 MB per GPU before the Ring traffic multiplier. The flat ring used one port's worth of fabric bandwidth because a single ring runs at its slowest edge; the hierarchical schedule uses all eight ports at once. That is also its precondition, since a node whose GPUs shared one uplink would carry the eight shard streams one after another and lose most of the speedup. This is why NVIDIA's NCCL and similar libraries often select hierarchical algorithms on multi-node clusters when the topology supports them.

This book's running instance makes the same comparison on the cluster it was defined for. Archetype A's 175B-parameter model produces a 350 GB BF16 gradient, and the reference cluster holds 1,024 GPUs as 128 nodes of 8. A single flat ring over every GPU would spend about 14.0 s per full-gradient reduction. The hierarchical schedule sends one-eighth of the gradient through each GPU's port in a shard ring that spans every node and completes in about 3.1 s, a 4.5× improvement. The cross-node phase still accounts for about 1.7 s of that time. Seconds per step remain a large exposed cost even after hierarchy. That is why the production layout for this instance shards the model so that each data-parallel group reduces a fraction of this payload. It is also why the overlap and compression techniques later in the chapter still matter at this scale.

These three phases also generalize to additional fabric tiers. For example, 4 racks of 4 nodes with 8 GPUs per node contain 128 GPUs. A rack-local ReduceScatter over 32 GPUs can leave each GPU with $1/32$ of the original payload for the cross-rack phase. The total elapsed time still includes both local phases, the cross-rack collective's traffic factor, startup, and contention. A cross-rack oversubscription ratio alone cannot determine a speedup over a flat ring.

5.7.2 Multidimensional reduction on a torus

Hierarchy fits a fabric whose tiers differ in speed. A torus, the direct-connect topology of Google's Tensor Processing Unit (TPU) pods, has uniform links instead, and its collectives reduce along one dimension at a time. TPU v4 pods support a 3D torus topology, unlike the 2D torus used by TPU v2

and v3. In the regular 3D torus model, each chip links directly to six neighbors $(\pm X, \pm Y, \pm Z)$ with equal per-link bandwidth. For collectives, the torus invites dimension-ordered reduction, which decomposes the AllReduce into sequential Ring AllReduces along axis-aligned torus rings. For a pod with $X \times Y \times Z$ chips, ranks first execute independent X-dimension rings for each fixed (Y, Z) coordinate, producing partial sums along each X line. A second phase runs Y-dimension rings to combine those partials across the $X \times Y$ plane. A third phase executes Z-dimension rings across fixed (X, Y) coordinates to complete the global reduction.

↩ The torus topology and its bisection and path-length trade-offs against the fat-tree are established in section 3.5.4.

On a 4×4 2D torus grid of 16 ranks, the dimension-ordered schedule runs in two phases (figure 5.8). In Phase 1 (orange solid paths), each row executes an independent 4-rank Ring AllReduce across the horizontal X dimension using dashed wrap-around boundary links to close the ring ($2(X-1) = 6$ steps). Once row partial sums are established, Phase 2 (slate solid paths) executes orthogonal 4-rank vertical Ring AllReduces down the Y dimension ($2(Y-1) = 6$ steps). These phases use direct nearest-neighbor links without multi-hop routing within a logical ring edge. They reduce the startup count from $2(N-1) = 30$ rounds on a flat 16-rank ring to $2(X-1) + 2(Y-1) = 12$ rounds, a 2.5× reduction in the modeled startup term. In this simple form each axis still processes the full message, so the round-count saving comes at a byte cost that the partitioned dimension-wise schedule removes.

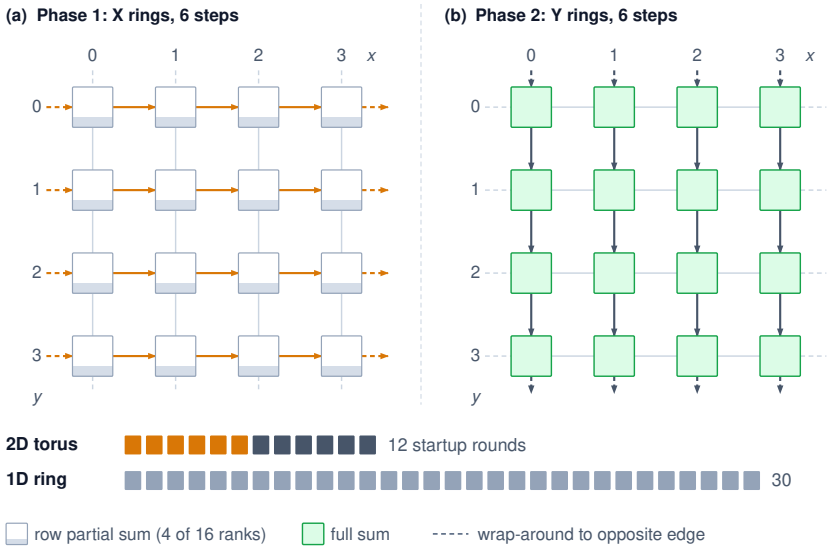

Figure 5.8: Dimension-Ordered Reduction on a Torus Topology: In a 2D torus mesh (4×4 node grid representing a simplified TPU pod slice), AllReduce proceeds sequentially along orthogonal dimensions. Phase 1 (orange) executes four parallel horizontal rings across the X dimension ($2(X-1) = 6$ steps). Phase 2 (slate) then executes four parallel vertical rings across the Y dimension ($2(Y-1) = 6$ steps) on the row-aggregated partial sums. Dashed stubs mark the torus wrap-around links, each leaving one edge of the grid and re-entering at the opposite edge. This dimension-ordered decomposition needs 12 startup rounds against 30 for a 16-rank 1D global ring, a $2.5\times$ reduction, but each axis processes the full message, so its bandwidth time is $3M/\beta$ for this grid.

The byte term does not shrink with the rings. Because every axis processes the full M-byte tensor, the bandwidth time is $2M[(X-1)/X + (Y-1)/Y + (Z-1)/Z]/\beta$ for equal per-axis bandwidth β. The startup term is $2[(X-1) + (Y-1) + (Z-1)]\alpha$. Shorter rings reduce startup cost, but these sequential full-message AllReduces do not retain the bandwidth cost of one global ring.

A bandwidth-efficient multidimensional implementation instead partitions the work with dimension-wise ReduceScatter followed by reverse-order AllGather. The X-dimension ReduceScatter sends $M(X-1)/X$ bytes per rank and leaves each rank at a given X coordinate holding the row-reduced sum for the same M/X shard. Because all ranks in a Y-column share that X coordinate, they hold matching shards and can run a Y-dimension ReduceScatter on that M/X slice alone, sending $(M/X)(Y-1)/Y$ bytes per rank and subdividing the slice to $M/(XY)$. On the 4×4 grid, the X-dimension ReduceScatter leaves each rank one of 4 shards of the tensor, and the Y-dimension

ReduceScatter leaves it one of 16, where the two axes finish the subdivision. On a three-axis torus, the Z-dimension ReduceScatter repeats the subdivision across ranks that share both X and Y coordinates, leaving each of the $N = XYZ$ chips with a distinct M/N shard of the global sum after sending $M(N-1)/N$ bytes in total. The mirrored AllGather along Z, Y, and X doubles that volume to $2(N-1)M/N$, the ring's optimal byte term, while keeping the short rings' startup count. The torus thus teaches the same lesson as hierarchy, which is to shrink the payload before each successive stage rather than send the full tensor through every one.

Dimension-ordered reduction follows direct neighbor links and separates axis traffic in time; it does not change the network diameter or eliminate contention with other workloads. The torus topology also provides natural fault tolerance through alternate routing paths. If one link in the X-dimension fails, traffic can detour through the Y or Z dimensions (at the cost of increased latency). The compiler and collective runtime must choose a schedule for the actual topology, message size, and available links.

5.7.3 In-network reduction with SHARP

Hierarchical and dimension-ordered schedules redistribute work across links, but their software reduction stages still require endpoint GPUs to receive and combine contributions. NVIDIA's Scalable Hierarchical Aggregation and Reduction Protocol (SHARP)[6] moves that intermediate aggregation into the network. Instead of routing each partial sum through a GPU for arithmetic, an aggregation-capable InfiniBand switch combines contributions as packets pass through it.

A centralized endpoint reduction and in-switch aggregation send different data up the switch tree (figure 5.9). In the endpoint tree of panel (a), intermediate switches merely forward raw packets, forcing all four full gradient tensors ($\mathbf{g}_0$ through $\mathbf{g}_3$) across the upper links into a single Reducer GPU that executes every addition. In panel (b), SHARP offloads arithmetic to aggregation-capable switch application-specific integrated circuits (ASICs): each intermediate switch computes partial sums ($\mathbf{g}_0 + \mathbf{g}_1$ and $\mathbf{g}_2 + \mathbf{g}_3$) in-flight, leaving only the combined shards for the root switch to finalize. In this four-GPU tree, moving reduction into the network fabric halves the upper link traffic and relieves accelerator endpoints from coordinating arithmetic steps.

6 **SHARP (Scalable Hierarchical Aggregation and Reduction Protocol)**: In-network computing on Quantum InfiniBand switches that performs reduction operations (such as sum, min, and max) as packets traverse an aggregation tree (Graham et al. 2020). The same work reports substantially higher reduction bandwidth than host-based algorithms and describes switch-resource limits on outstanding aggregation groups and trees, which can create contention when many collectives compete for the same in-network resources.

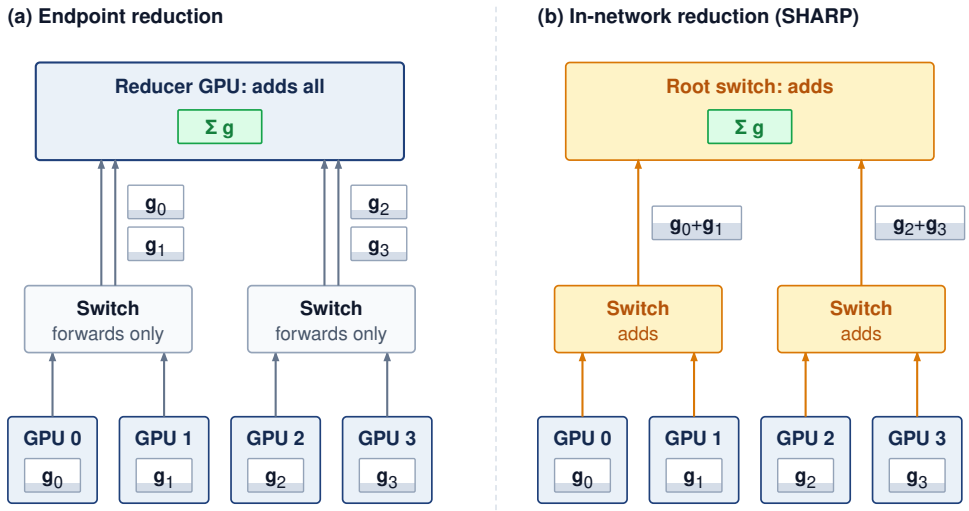

Figure 5.9: Endpoint and In-Network Reduction: The upstream half of a reduction over four GPUs. (a) Switches only forward packets, so every full gradient tensor crosses the upper links and converges on one reducer GPU, which does all the arithmetic. (b) Aggregation-capable switches add their inputs, so each upper link carries one partial sum and the root switch finishes the reduction; the result then travels back down the same tree. A software reduction tree also aggregates along its path, but performs arithmetic at endpoints rather than switches. The diagram compares where arithmetic happens, not measured speedup.

SHARP changes both terms of the cost model. On the startup side, it shortens the dependency chain. In a software-based tree reduction, each level waits for an intermediate GPU to receive,

reduce, and forward its inputs. With SHARP, the switch performs that step in the packet path, so no GPU execution stage or extra endpoint round trip sits between levels. On the byte side, each upper-tier link carries one partial sum rather than every contribution beneath it (figure 5.9). In the idealized aggregation tree, each GPU sends its contribution up once and receives the result once, whereas a ring rank sends nearly $2M$. Endpoint memory traffic falls for the same reason. Switch processing and contention for aggregation resources add costs of their own, so moving arithmetic into the network changes the α-β terms rather than eliminating them.

Which benefit matters follows the regimes of section 5.5. For small and medium messages, where startup rounds dominate, the shorter dependency chain is the main gain, although for very small messages fixed switch-processing overhead can limit it. For large messages the startup saving becomes proportionally small, and what remains is the byte saving at the endpoints and on the upper tier, which shortens the collective only when those links are the limiting resource. Graham et al. (2020) report multi-fold reduction-bandwidth improvements and smaller application-level gains for PyTorch workloads when SHARP Streaming-Aggregation, its mode for large messages, is enabled on HDR/Quantum InfiniBand switches.

A faster collective reaches step time only through the exposed part of its transfer, and strong scaling campaigns (training a fixed model across more accelerators without proportionally growing batch size) enlarge that part. As the per-accelerator batch shrinks, each gradient tensor stays large (proportional to model size) but the backward pass per accelerator shortens, exposing more of the communication budget. The limiting term can still be bandwidth; strong scaling alone does not establish that startup latency dominates. In-switch aggregation helps when its reduction in endpoint work or fabric traffic shortens the exposed critical path.

In-switch reduction also brings hardware and capacity constraints of its own. The switch ASIC must support the exact reduction operator and numerical format the job uses, such as FP16 or BF16 gradients, which vary across switch generations. Its finite internal memory limits how many aggregation trees can remain active at once. A large cluster running many concurrent jobs or process groups can therefore exhaust switch-tree capacity and fall back to host-based reductions. Moreover, the HDR and Quantum SHARP implementations examined here depend on InfiniBand switch support, so their gains do not carry over to Ethernet fabrics unless an equivalent in-network aggregation path is present.

5.7.4 Topology-aware routing

Two runs of the same 64-GPU Ring AllReduce on identical hardware can differ by a multiple in throughput depending on whether consecutive logical ranks sit on the same NVLink domain or across shared inter-node uplinks. Hierarchy and in-network aggregation change where bytes move and where sums are computed, but their benefit still depends on mapping the logical communication pattern onto physical links.

Communication libraries like NCCL perform topology detection at initialization and use graph search to choose paths through the hardware topology available to them (NVIDIA 2026c). For each pair of ranks, the library must determine whether the peer sits behind a local NVLink switch or 100 meters away across an InfiniBand fabric, because that mapping decides which physical links each logical hop occupies. Local GPU, NVLink, PCIe, and NIC connectivity informs those choices, but local discovery does not reveal every remote switch, route, or competing flow. Fabric configuration and placement information must complement the library's view when diagnosing cross-node contention.

The topology detection process begins with hardware enumeration. NCCL queries the PCIe bus to discover GPU placement, NVLink connectivity between GPUs, and NIC-to-PCIe-switch affinity. From this information, it constructs an internal graph where nodes represent GPUs and edges represent physical links with annotated bandwidth and latency. A graph search algorithm then finds communication paths that maximize aggregate bandwidth while minimizing the number of cross-domain hops (transitions between NVLink, PCIe, and InfiniBand domains). This topology-aware path selection is what allows NCCL to approach theoretical peak bandwidth on well-configured systems, while topology-unaware implementations can leave a large fraction of bandwidth unused.

5.7.4.1 Rail-optimized routing (NVIDIA DGX)

In the eight-GPU DGX layout considered here, each GPU has a dedicated inter-node NIC, and the fabric connects matching NIC positions to rail switches. Rail-optimized routing fits the corresponding-shard phase of hierarchical AllReduce. GPU 0 on Node A exchanges its shard with GPU 0 on Nodes B and C, GPU 1 exchanges with peer GPU 1s, and so on. The eight rails can carry these groups concurrently (table 5.8), subject to their physical capacity and competing traffic. Other collectives may require cross-rail paths.

↩ Rail-optimized cabling, which connects each GPU slot position to its own rail switch, is established in section 3.5.3.

Table 5.8: Rail-Optimized Traffic Distribution: Each rail carries 1/8 of the total traffic independently.

Rail	Participants	Traffic
Rail 0	Node0-GPU0 ↔ Node1-GPU0 ↔ Node2-GPU0 ↔ …	$M/8$ each
Rail 1	Node0-GPU1 ↔ Node1-GPU1 ↔ Node2-GPU1 ↔ …	$M/8$ each
…	…	…
Rail 7	Node0-GPU7 ↔ Node1-GPU7 ↔ Node2-GPU7 ↔ …	$M/8$ each

If all eight shard streams converge on one remote NIC, that endpoint faces 8× the balanced receive demand. In the equal-shard layout, matching each shard group to its intended rail distributes one-eighth of the payload to each NIC. This removes that particular concentration; achieving the predicted rate still requires sufficient switch, uplink, and bisection capacity for concurrent traffic.

NVIDIA designed the DGX node and rail layout together so that Phase 2 of hierarchical AllReduce (section 5.7.1) maps each corresponding-rank shard group onto its own rail. NCCL exploits that structure whenever its topology detection finds aligned ranks. When the job scheduler assigns process ranks without topology awareness, however, logical shards and physical rails drift apart. Cross-node traffic is then routed through the wrong NICs or across spine switches, cutting a node's usable injection bandwidth to as little as one-eighth of the rail-aligned value when all eight shard exchanges converge on one NIC. Large deployments therefore enforce topology-aware rank assignment through device ordering (`CUDA_VISIBLE_DEVICES`), NIC binding, and scheduler placement policies so that logical rank i on every node stays bound to rail i.

Hierarchical execution and topology-aware routing connect the cost model to a physical placement. The α-β terms estimate the service required by each scheduled exchange; hierarchy redistributes work across fabric tiers, and rail alignment can balance corresponding-shard traffic across independent ports. None removes shared-capacity limits. Measured collective duration must confirm that the intended placement shortens the dependency path before the design is accepted.

↪ How the scheduler places a job's gangs on the topology hierarchy so that its collectives stay on fast links is developed in section 8.6.

5.8 AllToAll Exchanges at Scale

Hierarchy, in-network reduction, and rail alignment were all developed for AllReduce, whose uniform payload lets every rank follow the same schedule. In a mixture-of-experts model with 64 experts distributed across 8 GPUs (8 experts per GPU), or in the sharded embedding tables of table 5.1, a gating network or lookup table assigns each input to a specific destination rank. Tokens assigned to remote GPUs must physically cross the fabric before expert computation can proceed and return afterward, breaking AllReduce's symmetry so the exchange finishes only when its slowest path completes.

> **Systems Perspective 5.1: AllToAll vs. AllReduce: Why scale differs**
>
> Ring AllReduce can pipeline a uniform payload through neighbor exchanges. AllToAll presents a different scaling problem because every rank has destination-specific traffic, but the logical peer count alone does not determine its cost.
>
> An AllToAll over N ranks creates $N(N-1)$ logical source-destination transfers, yet each rank's aggregate payload can remain fixed under weak scaling (per-rank work held fixed as ranks are added). Its lower bound is governed by the largest of four terms: source injection time (how long a rank's port takes to send its bytes), destination reception time (how long a port takes

to receive them), traffic crossing the fabric bisection, and per-message startup or scheduling overhead. Oversubscription and routing determine how much physical capacity those flows share. Message granularity determines how often they pay α. In the α-β model, that startup term acts as a bound only when messages are serialized, because concurrent sends can share one startup, although the rate at which an endpoint can issue messages still limits those sends. Routing skew determines whether one destination or link becomes the critical path. A full-bisection network carrying balanced, coarse-grained traffic can therefore scale well despite the quadratic logical connection count, whereas an oversubscribed fabric or skewed expert assignment can stall at much smaller N.

Expert parallelism and large-scale recommendation systems can reach a communication wall earlier than standard data parallelism when their traffic exhausts injection, reception, or bisection capacity, or when routing and destination skew create hotspots. The primitive determines the traffic matrix; the physical fabric and routing distribution determine the scaling ceiling.

Pricing that 64-expert, 8-GPU MoE layer under uniform routing establishes the baseline transfer volume and injection bound before dynamic skew enters.

Napkin Math 5.5: AllToAll for MoE token routing

Problem: An MoE model processes a batch of 4096 tokens across 8 GPUs (512 tokens per GPU). Each token is a 2048-dimensional hidden state in BF16 (4.096 KB per token). The gating network assigns each token to exactly 1 of 64 experts (8 experts per GPU). Assuming uniform routing (each expert receives 64 tokens), how much data does each GPU send and receive?

Math:

Each GPU holds 512 tokens that need to reach 64 different experts across 8 GPUs. With uniform routing, each GPU sends 64 tokens to each of the other 7 GPUs (keeping 64 tokens local).

- Data per GPU-to-GPU transfer: 64 tokens $\times$ 4.096 KB/token = 262.144 KB
- Total data sent per GPU: 7 $\times$ 262.144 KB = 1.84 MB
- Total data received per GPU: 1.84 MB (symmetric)

Latency Analysis (InfiniBand NDR, $\alpha = 1.5$ µs, $\beta = 50$ GB/s):

Each of the 7 transfers is a 262.144 KB message. From the α-β model: T_{transfer} = 1.5 µs + 5 µs = 6.5 µs per transfer.

If serialized: 7 $\times$ 6.5 µs = 45.5 µs. If all 7 transfers are issued concurrently, they still share the GPU's 50 GB/s injection port. The aggregate injection lower bound is therefore $\alpha +$ total bytes sent$/\beta = 38.2$ µs, before contention and additional collective overhead. Full duplex permits simultaneous sending and receiving; it does not provide one full-bandwidth outgoing port per peer.

Systems insight: The injection result is only a lower bound. If the topology's usable bisection bandwidth is $\text{BW}_{\text{bisect}}$ and V_{cut} bytes must cross its minimum cut, the exchange also obeys $T \geq V_{\text{cut}}/\text{BW}_{\text{bisect}}$. Oversubscription reduces $\text{BW}_{\text{bisect}}$, fine-grained peer messages add startup cost, and expert skew raises the maximum receive volume above the uniform estimate. The forward pass runs the AllToAll twice per layer, once for token dispatch and once for result collection; training also communicates gradients during backward.

5.8.1 Dynamic routing and expert imbalance in MoE AllToAll

The worked trace in notebook 5.5 assumes uniform token routing, in which every expert receives an identical share of tokens. In production MoE models, routing decisions are dynamic and data-dependent. The exchange keeps its two-phase shape, but its buffer sizes, its slowest receiver, and its load on the fabric all become functions of the routing distribution. The hierarchy just developed for AllReduce reappears as one of the schedules that can contain them.

Each MoE layer performs two exchanges. In an MoE transformer layer with N_{experts} experts distributed across N ranks (N_{experts}/N experts per rank), each rank holds a local batch of B_{tok} tokens. A gating router computes top-k_{experts} routing probabilities over all N_{experts} experts for each token, inducing two collective communication phases separated by expert computation. During token dispatch, each rank partitions its B_{tok} local token activation vectors ($\mathbf{x} \in \mathbb{R}^d$, where d is the model's hidden dimension) according to the destination ranks hosting the selected experts. The rank then executes an AllToAll collective (or variable-size `AllToAllv`) to route tokens across the fabric according to the routing matrix. After each rank computes forward passes through its assigned local expert feed-forward networks (FFNs), the token combine phase executes a second AllToAll collective to route output vectors back to their originating ranks. Each rank then computes the weighted combination of expert outputs:

$$\mathbf{y}_i = \sum_{j \in \text{top-}k_{\text{experts}}} p_{\text{gate}_{i,j}} \cdot \text{Expert}_j(\mathbf{x}_i)$$

Because routing is data-dependent, so are the required buffer sizes, even when the input tensor shape is fixed. One simplified model treats assignments of different tokens independently: the count for expert e is then $B_e \sim \text{Binomial}(NB_{\text{tok}}, p_e)$, where p_e is the fixed probability that a token selects expert e. Correlated tokens or changing router probabilities require a different count model. To avoid reallocating buffers for every routing decision, a capacity-limited implementation can reserve space using an expert capacity factor $C_{\text{cap}} \geq 1.0$, a multiplier on each expert's average share of token assignments (Lepikhin et al. 2021; Fedus et al. 2022):

$$\text{Capacity}_e = \left\lceil \frac{N \times B_{\text{tok}} \times k_{\text{experts}}}{N_{\text{experts}}} \times C_{\text{cap}} \right\rceil$$

The capacity factor exposes a trade-off between reserved space and retained expert assignments. At uniform expected capacity ($C_{\text{cap}} = 1.0$), some experts can still overflow while others have unused slots. A dropping implementation bypasses excess assignments through the residual connection, changing the computation and potentially harming training quality. Increasing headroom ($C_{\text{cap}} > 1.0$, for example 1.25–1.5) can reduce overflow. If the implementation transmits every reserved slot, its padded payload grows by approximately a factor of C_{cap} relative to the uniform capacity-one budget; kernels that skip unused slots need not perform the corresponding dummy computation. A dropless implementation retains every assignment and can use variable-size `AllToAllv` to avoid sending unused slots. It still needs enough memory for the realized counts and remains exposed to routing skew.

Because token routing is determined by input semantics, some experts inevitably receive disproportionate token volumes. An expert that receives several times the average load raises its host GPU's receive traffic by the same excess. The host's total receive volume depends on the combined load of all experts assigned to it, which is the rank-level quantity in the bound that follows.

Routing skew can create stragglers in both communication and expert computation. An inbound reception bottleneck bounds dispatch completion by the most heavily loaded receiver. Other ranks may finish their transfers earlier, but work that needs the delayed expert outputs must wait. The receive budget gives a lower bound before expert computation and the return exchange are counted:

$$T_{\text{dispatch}} \geq \frac{B_{\text{max}} \times d \times s_{\text{elem}}}{\text{BW}_{\text{rx}}}$$

where $B_{\text{max}} = \max_i \sum_{e \in \text{rank } i} B_e$ counts the largest remotely received token batch at any rank i, restricting B_e in this sum to assignments arriving from other ranks. Here s_{elem} is the bytes per activation element, and BW_{rx} is that rank's receive bandwidth.

Bisection capacity provides a separate limit. Ring AllReduce exchanges data with logical ring neighbors, whereas AllToAll permits every rank on one side of the bisection to send to every rank on the other. Under uniform destination assignment, a fraction $(N-1)/N \approx 1$ of tokens leaves its source rank; the fraction crossing nodes or a particular fabric cut depends on placement. In the simplified symmetric fabric model, each rank has injection bandwidth $\text{BW}_{\text{inject}}$, the rate its port can send into the fabric, and an oversubscription ratio $R : 1$ gives a switch tier R times more downlink

than uplink capacity. Usable bisection bandwidth $\text{BW}_{\text{bisect}} = (N \cdot \text{BW}_{\text{inject}})/(2R)$ then bounds the time needed to carry traffic across that cut:

$$T_{\text{dispatch}} \geq \frac{V_{\text{cut}}}{\text{BW}_{\text{bisect}}}$$

↩ Bisection bandwidth and the cost of oversubscription are established in section 3.5.

Napkin Math 5.6: The cost of a hot expert rank

Problem: Return to the uniform routing of notebook 5.5: 8 GPUs, 64 experts, 4.096 KB per token, and one 50 GB/s port per GPU. Suppose skew sends one rank's experts 1.5 times the average remote load. This is an assumed realized routing skew; the per-expert capacity factor reserves total-assignment headroom and does not determine remote arrivals. How does the receive bound on dispatch change?

Math:

1. **Uniform receiver**: Each GPU receives 448 remote tokens, or 1.84 MB, so $T_{\text{dispatch}} \geq 1.84$ MB / 50 GB/s ≈ 36.7 µs.
2. **Hot receiver**: B_{max} rises to 672 remote tokens, or 2.75 MB, so $T_{\text{dispatch}} \geq 2.75$ MB / 50 GB/s ≈ 55.1 µs.

Systems insight: The exchange cannot finish before its hottest receiver, so a 1.5× load on one rank raises the dispatch bound for the whole layer by the same factor. Transmitting unused capacity slots can add padding bytes; the amount depends on the packing layout rather than on this receive-skew factor alone. A variable-size exchange removes unused-slot traffic but not the receive bound. Expert computation processes local as well as remote assignments, so its work grows by the same factor only if the hot rank's total local-plus-remote assignment count also grows by that factor, with unchanged work per assignment.

Scale compounds both problems. To amortize small-message overhead across thousands of GPUs, hierarchical MoE runtimes decompose the flat AllToAll into a two-level collective (Hwang et al. 2023). Intra-node NVLink first coalesces tokens destined for experts on the same remote server node. Participating nodes then exchange aggregated payloads over InfiniBand or RoCE, and receiving GPUs use NVLink to distribute tokens to the GPUs hosting the targeted experts. At the logical level, this replaces $N(N-1)$ directed rank pairs with $N_{\text{nodes}}(N_{\text{nodes}}-1)$ directed node pairs; actual transport connections depend on the implementation. Coalescing can improve link utilization and reduce startup overhead, but the token bytes still have to cross the required fabric cuts, and skew can still overload a destination.

5.8.2 AllToAll schedule selection

A flat, fixed-size AllToAll schedule that works well for balanced, megabyte-scale peer transfers on eight GPUs can stall across thousands of GPUs when per-peer chunks shrink to a few kilobytes or when expert skew forces every rank to pad its buffers to the hottest receiver's count. Because `AllToAll` names only the destination-specific contract, the runtime must match its schedule to the traffic matrix and physical hierarchy rather than treating every exchange as one flat pattern. Three schedules cover the common regimes (table 5.9). Direct pairwise exchange suits balanced, coarse per-peer chunks. Hierarchical exchange coalesces many small remote chunks before they cross scarce inter-node links. Variable-size exchange avoids padding when token or embedding counts differ by destination. Hierarchical designs such as Tutel's two-dimensional exchange demonstrate why the best schedule changes with message size, scale, and topology (Hwang et al. 2023).

Table 5.9: AllToAll Schedule Selection: The right schedule minimizes the active bound rather than the logical peer count. Direct exchange preserves simplicity, hierarchy coalesces traffic before scarce links, and variable-size exchange removes padding without removing skew.

Schedule	When it fits	Cost exposed
Direct pairwise exchange	Balanced, coarse per-peer chunks on a small or full-bisection process group	Each rank injects directly to its peers; fine granularity repeats startup and scheduling cost, while oversubscription can bind the bisection
Hierarchical/coalesced exchange	Many small remote chunks across multi-GPU nodes or rails; pack by remote node, exchange aggregates, then redistribute locally	Coalescing amortizes inter-node startups and respects fast local links, but adds packing, local copies, synchronization, and temporary buffers
Variable-size exchange	Token or embedding counts differ by destination and padding to the maximum count would waste substantial bandwidth	Explicit counts and displacements avoid padding, but count exchange and irregular transfers add control overhead; the hottest receiver still sets completion time

The selection test follows the four terms from perspective 5.1 and section 5.8.1. When per-peer chunks are already coarse and source injection or fabric bisection dominates, coalescing cannot shrink the byte term and only adds local copies. When each remote destination receives a small chunk so that the serialized startup cost $N\alpha$ dominates instead, a hierarchical schedule replaces those many inter-node startups with a single aggregated exchange per remote node. Under routing skew, a variable-size interface keeps padding from inflating V_{cut}, though the runtime must still provision memory and time for the hottest receiver's load. Because these bounds depend on live routing behavior, a runtime must benchmark candidate schedules under the observed token-count distribution, rank placement, and fabric topology rather than relying on a balanced, fixed-size microbenchmark. Even the best schedule, however, must still transmit the full activation or gradient payload unless the system reduces the bytes themselves.

5.9 Gradient Compression Under Bandwidth Scarcity

Training across data centers over wide-area links or on cloud instances with commodity Ethernet can leave each worker with 10–100× less network bandwidth than the reference InfiniBand fabric. Federated training on phone-class devices, the Archetype C (Federated MobileNet) regime, faces cellular or Wi-Fi uplinks that are both slower and intermittently available. Even after hierarchical routing and schedule tuning extract all usable capacity from such links, moving a full-precision gradient or a multi-step client update can still dominate step time. Gradient compression attacks the payload size directly by trading numerical representation fidelity for fewer transmitted bits, offering nominal value-payload reductions of 4–1000× before metadata and aggregation costs reduce the realized savings. Because noisy or biased updates can slow or stall optimization, the central engineering question is how much communication volume a compressor can remove while preserving progress toward the training quality target.

↪ The per-round uplink budget of federated training, and the update compression that fits it to cellular and Wi-Fi links, are developed in section 11.7.2.

5.9.1 Precision reduction via quantization

Quantization is the least disruptive compression lever when bandwidth is the limiting resource, because it keeps every gradient element but reduces the bits used to represent it. Most gradients are computed in FP32 (32-bit floating point) or BF16 (16-bit brain float), so reducing bit-width directly reduces communication volume. The useful question is how much gradient fidelity the optimizer can absorb across the progression from mild to aggressive quantization. The compression ratios discussed in this section follow the compression literature and are quoted against FP32 values. The chapter's 70B payload is already BF16, so against it every ratio halves; INT8 saves 2× rather than 4×, and 1-bit encoding 16× rather than 32×.

FP16 halves FP32's value bytes when its range supports the workload, and INT8 stores each quantized value in one byte, a $4\times$ reduction relative to FP32 before scale metadata. A symmetric INT8 scheme uses 255 levels from negative 127 through positive 127. For a nonzero gradient vector, it computes a tensor-wide scaling factor $s = \max(|\mathbf{g}|)/127$, quantizes via $\mathbf{g}_{\text{int8}} = \text{round}(\mathbf{g}/s)$, and has the receiver reconstruct $\hat{\mathbf{g}} = \mathbf{g}_{\text{int8}} \times s$. An all-zero vector requires a separate zero case.

Unlike quantized weights, quantized gradients pass through a reduction that combines them across ranks. Each rank computes its scale from its own gradient, so independently scaled integer vectors cannot be added directly to recover the sum of their reconstructed gradients. The collective must carry every rank's scale through the reduction and use sufficient accumulator precision, or reconstruct values before reducing them.

At the extreme, 1-bit SGD transmits only the sign of each gradient element ($+1$ or -1), yielding a $32\times$ reduction in value bits over FP32. The receiver reconstructs values using a learned or adaptive scaling factor, though the substantial noise introduced by 1-bit quantization can severely degrade convergence without error-feedback mechanisms (Seide et al. 2014; Karimireddy et al. 2019).

↪ Byte widths for every format named here are collected for reference in section A.6.

5.9.1.1 Block quantization for gradient communication

Using a single scaling factor for an entire tensor makes every element's quantization resolution depend on the largest magnitude. When magnitudes vary substantially across the tensor, small gradients may round to zero even in regions whose values could be represented accurately with a smaller scale. Block quantization addresses this by dividing the gradient tensor into blocks of b_{block} elements (for example, $b_{\text{block}} = 64$ or 128) and computing an independent scaling factor per block. Each block's scaling factor adapts to its local range, improving resolution where that range is smaller at the cost of transmitting $d_{\text{grad}}/b_{\text{block}}$ additional scaling factors for evenly divided blocks.

For a gradient vector of dimension d_{grad} quantized to INT8 with block size b_{block}, the message size is $d_{\text{grad}} \times 1$ byte for the quantized values plus $(d_{\text{grad}}/b_{\text{block}}) \times 4$ bytes for the FP32 scaling factors, or $d_{\text{grad}} + 4d_{\text{grad}}/b_{\text{block}}$ bytes in total. The effective compression is therefore $4d_{\text{grad}}/(d_{\text{grad}} + 4d_{\text{grad}}/b_{\text{block}}) = 4b_{\text{block}}/(b_{\text{block}} + 4)$, which for $b_{\text{block}} = 128$ reaches $3.88\times$, close to the theoretical $4\times$. For rounding to the nearest level without clipping, the absolute error bound changes from $s \cdot 0.5$, where s is the global scaling factor, to $s_b \cdot 0.5$, where s_b is the block-local scaling factor. Blocks with smaller ranges obtain tighter bounds; the improvement depends on how magnitudes are distributed among blocks.

Interfaces that fuse quantization and reduction must preserve the meaning of each block's scale throughout aggregation. They must also control the resulting numerical error. Deterministic rounding generally introduces bias, whereas suitable randomized quantizers preserve the input in expectation while adding variance. Error feedback addresses accumulated compression error, but the need for it depends on the compressor and optimizer rather than on bit-width alone. Quantization should therefore move only as far down the precision ladder as the convergence budget permits.

5.9.2 Gradient sparsification methods

Even 1-bit quantization must send at least one bit per coordinate, capping its value-payload reduction at $16\times$ relative to BF16 ($32\times$ relative to FP32) before scale metadata. Pushing compression to $100\times$ or beyond requires an orthogonal lever, sparsification, which keeps full or half precision for selected coordinates and transmits only a small subset of gradient elements while setting the rest to zero.

5.9.2.1 Top-k sparsification

Top-k compression selects that subset by magnitude. For a gradient vector $\mathbf{g} \in \mathbb{R}^{d_{\text{grad}}}$, it transmits only the K_{top} elements with the largest absolute magnitude, setting the rest to zero (Aji and Heafield 2017; Lin et al. 2018):

$$\text{TopK}(\mathbf{g}) = \mathbf{g} \odot \mathbf{1}_{|\mathbf{g}| \geq |\mathbf{g}|_{(K_{\text{top}})}}$$

where $|\mathbf{g}|_{(K_{\text{top}})}$ is the K_{top}-th largest element by magnitude. The threshold expression retains all ties; an implementation targeting an exact count must break ties at the boundary. With $K_{\text{top}} = 0.001 \times d_{\text{grad}}$ (keeping only 0.1 percent of elements), the retained values occupy $1000\times$ less space before index overhead.

The compression headline must be discounted by the cost of encoding the sparse representation. Transmitting a sparse gradient requires sending both the nonzero values and their indices. For a gradient vector of dimension d_{grad} with K_{top} nonzero elements, the encoded message contains K_{top} values (each in FP32 or FP16) plus K_{top} indices (typically INT32). The total message size is

$K_{\text{top}} \times (4+4) = 8K_{\text{top}}$ bytes for FP32 values with INT32 indices. The effective compression ratio is therefore $4d_{\text{grad}}/(8K_{\text{top}}) = d_{\text{grad}}/(2K_{\text{top}})$.

For $K_{\text{top}} = 0.001d_{\text{grad}}$, the encoded message yields a compression ratio of 500×, not the 1000× suggested by the kept-element fraction alone, because the index overhead is significant. Using a 16-bit unsigned index representation (for $d_{\text{grad}} < 65536$) or run-length encoding for suitable sparsity patterns can improve the effective ratio. Aggregation imposes a separate cost. Workers may retain different coordinates, so summing their packed value arrays would combine unrelated parameters. A sparse collective must match indices and sum contributions to each coordinate; the union of selected coordinates can be much larger than any one worker's selection. Local compression ratios therefore do not directly specify collective traffic or runtime.

Top-k is not the only sparsification rule. Random-k sparsification selects K_{rand} elements uniformly at random and scales them by $d_{\text{grad}}/K_{\text{rand}}$ to maintain an unbiased gradient estimate. Random-k has the advantage of being an unbiased compressor even without error feedback, because $\mathbb{E}[\text{RandomK}(\mathbf{g})] = \mathbf{g}$. Its random selection can discard large components, and rescaling the retained values introduces variance. Top-k retains the largest components but introduces bias. Their convergence behavior therefore depends on the optimizer, residual correction, and gradient distribution; unbiasedness alone does not determine which reaches a quality target sooner.

5.9.2.2 The convergence problem

Naively discarding small gradients with Top-k creates a systematic bias. If a parameter's gradient remains below the selection threshold, even a consistent contribution of 0.01 per step is repeatedly omitted. Over thousands of steps, those omissions can materially change the optimization trajectory. The problem is persistent exclusion, which can occur even when each discarded component looks negligible in isolation.

The bias of uncorrected Top-k means that the usual argument for SGD with an unbiased gradient estimate no longer applies directly. Biased compression can cause convergence failure even when the uncompressed method converges. Deterministic quantization can also repeatedly omit small contributions; randomized quantization and rescaled Random-k instead require accounting for their added variance. Error feedback addresses accumulated compression error for quantized gradients (Seide et al. 2014) and for sparsified gradients (Stich et al. 2018; Karimireddy et al. 2019).

5.9.3 Error feedback and residual accumulation

Error feedback prevents that 0.01 sub-threshold component from being lost forever. Rather than discarding unselected or rounded-away gradient coordinates, each worker retains a residual accumulator $\mathbf{e}_t$ in local accelerator memory. At step t, the system adds $\mathbf{e}_t$ to the new gradient $\mathbf{g}_t$, compresses the corrected vector, transmits $\mathbf{v}_t$, and stores the untransmitted difference as $\mathbf{e}_{t+1}$, allowing persistent small contributions to accumulate until they cross the transmission threshold.

Definition 5.3: Error feedback mechanism

Error feedback is a distributed-training technique for gradient compression that maintains a per-worker residual accumulator $\mathbf{e}_t$, re-injecting the compression error back into the next gradient update so that information omitted from the transmitted update remains in the residual. A finite run can end with an untransmitted residual.

1. **Significance**: Top-k sparsification at 1 percent keeps only the largest 1 percent of gradient values by magnitude, reducing each worker's retained value payload from 140 GB to 1.4 GB for a full 70B gradient, a 100× reduction before sparse-index and aggregation overhead. Without error feedback, the discarded 99 percent of gradient values create a biased update that can harm convergence. With error feedback, omitted information remains in the residual and can influence later selections, supporting the convergence guarantees analyzed for memory/error-feedback variants under their assumptions (Stich et al. 2018; Karimireddy et al. 2019).
2. **Distinction**: Unlike plain lossy compression, which permanently discards sub-threshold gradient components, error feedback delays their transmission. A component too small

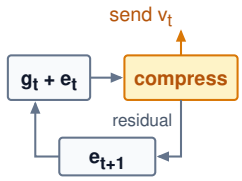

Error feedback carries the residual into the next compressed step.

to send in one step stays in the residual and can influence later transmitted updates as it accumulates.

3. **Common pitfall**: A frequent misconception is that error feedback is automatic whenever a compressor is used. It is not: the training system must explicitly preserve and reinject the residual. Without that residual path, Top-k and 1-bit quantization introduce systematic bias that can accumulate across steps (Karimireddy et al. 2019).

Error feedback preserves cumulative gradient information over time. Consider what happens over K steps:

$$\sum_{t=1}^{K} \mathbf{v}_t = \sum_{t=1}^{K} [(\mathbf{g}_t + \mathbf{e}_t) - \mathbf{e}_{t+1}] = \sum_{t=1}^{K} \mathbf{g}_t + \mathbf{e}_1 - \mathbf{e}_{K+1}$$

If the error accumulator remains bounded under the chosen compressor and training assumptions, then as $K \to \infty$:

$$\frac{1}{K} \sum_{t=1}^{K} \mathbf{v}_t \to \frac{1}{K} \sum_{t=1}^{K} \mathbf{g}_t$$

Under that bounded-residual condition, the difference between the two averages vanishes. At a finite stopping point, untransmitted information remains in the residual. Repeated small gradients can accumulate in $\mathbf{e}_t$ until they exceed the compression threshold, as the worked trace illustrates.

The telescoping property of compression error across time is why error feedback can transform a biased compressor into a convergent training method. Theoretical analyses show convergence guarantees for sparsified SGD with memory and for EF-SGD with broad compression operators under their stated assumptions (Stich et al. 2018; Karimireddy et al. 2019). A rounding compressor with no residual path shows what is lost without that memory (table 5.10). It transmits each value rounded to the nearest whole unit, so any value below 0.5 is sent as zero.

Table 5.10: Naive Compression Trace (No Error Feedback): Five-step worked trace showing how a rounding compressor transmits zero whenever an individual gradient falls below 0.5.

Step	True gradient g_t	Transmitted v_t	Cumulative transmitted	Cumulative true
1	0.4	0	0	0.4
2	0.3	0	0	0.7
3	0.2	0	0	0.9
4	0.4	0	0	1.3
5	0.3	0	0	1.6

A residual accumulator $\mathbf{e}_t$ that re-injects the discarded remainder into the next step changes the outcome of the same sequence (table 5.11). Instead of vanishing at each rounding step, the fractional updates accumulate locally until they cross the transmission threshold.

Table 5.11: Error-Feedback Compression Trace: Same five-step sequence under error-feedback compression.

Step	g_t	e_t	$g_t + e_t$	v_t	$e_{t+1} = (g_t + e_t) - v_t$	Cumulative v
1	0.4	0	0.4	0	0.4	0
2	0.3	0.4	0.7	1	–0.3	1
3	0.2	–0.3	–0.1	0	–0.1	1
4	0.4	–0.1	0.3	0	0.3	1
5	0.3	0.3	0.6	1	–0.4	2

With error feedback, the five-step run transmits 2 and ends holding –0.4 in the residual. Transmitted plus residual, 2 + (–0.4), equals the true total of 1.6, whereas rounding alone transmitted 0. The residual conserves the total, but convergence still depends on the compressor and optimizer assumptions.

5.9.4 1-bit Adam compression-aware optimization

Compressing raw gradients is easiest to control under stochastic gradient descent, where the parameter update is linear in the gradient. Adaptive optimizers such as Adaptive Moment Estimation (Adam) divide the first moment (momentum) by the square root of a running second-moment (variance) estimate, so compression noise in the raw gradient distorts that nonlinear denominator and can destabilize training. Microsoft's DeepSpeed team observed that Adam's second-moment estimate stabilizes early in training even while momentum continues to evolve (Tang et al. 2021). Rather than compressing raw gradients on every step, 1-bit Adam runs standard full-precision Adam during an initial warmup phase until the variance estimate settles. It then freezes that variance tensor as a fixed adaptive preconditioner and switches to a communication-efficient phase. In this phase, workers communicate only the momentum using 1-bit compression (sign plus a scaling factor) while retaining each compression residual locally for the next step.

Because the frozen preconditioner is no longer perturbed by quantization noise and the local residual feeds omitted momentum back into subsequent messages, the method controls accumulated compression error and preserves Adam's convergence trajectory without requiring each 1-bit message to be an unbiased estimator. Across experiments up to 256 GPUs, Tang et al. (2021) report up to 5× end-to-end communication-volume reduction. That is smaller than the raw 32× ratio of 1-bit values to FP32 momentum, because the uncompressed warmup and other communication costs are included. The same work reports convergence speed that matches uncompressed Adam, with up to 3.3× higher throughput for BERT-Large pretraining and up to 2.9× higher throughput for Stanford Question Answering Dataset (SQuAD) fine-tuning. Designing the compressor around the optimizer's internal state thus succeeds where replacing every raw gradient with its sign would fail, provided the compression overhead remains small compared with the communication time saved.

5.9.5 Bandwidth versus convergence trade-offs in compression

Across quantization, sparsification, error feedback, and 1-bit Adam, greater byte reduction demands stronger residual or preconditioner safeguards to protect convergence. Placing these methods side by side against both FP32 and BF16 baselines (table 5.12) maps each compression ratio to the convergence control and link regime it requires (Seide et al. 2014; Aji and Heafield 2017; Lin et al. 2018; Stich et al. 2018; Karimireddy et al. 2019; Tang et al. 2021).

Table 5.12: Gradient Compression Methods: Encoded-payload ratios relative to an FP32 payload, as the literature quotes them, and relative to the BF16 payload this chapter's examples use. Top-k assumes FP32 values and 32-bit indices in both columns; quantization scales and other metadata reduce the stated upper bounds. No method has a fixed convergence penalty across workloads. The right choice depends on whether bandwidth or convergence speed is the binding constraint; evaluate the complete time to the same quality target, including encoding and any additional training steps.

Method	Ratio vs. FP32	Ratio vs. BF16	Convergence impact	Best use case
FP16	2×	1× (none)	Check range and numerical stability	Common accelerator baseline
INT8 + Error FB	At most 4×	At most 2×	Workload- and scaling-dependent	Bandwidth-constrained clusters
Top-k (1%) + Error FB	At most 50×	At most 25×	Depends on sparsity and residual control	Cross-data-center training
1-bit + Error FB	At most 32×	At most 16×	Requires compression-aware optimization	Extreme bandwidth constraints

5.9.5.1 Decision criteria for gradient compression

The decision rule is not the compression ratio alone. Compression is worthwhile only when the wall-clock time it saves across the run exceeds its encoding overhead plus the cost of any extra steps

that slower convergence requires, so the comparison must be made in time to the same quality target. A high communication-computation ratio $\rho = T_{\text{comm}}(N)/(T_{\text{compute}}/N)$ creates an opportunity when that communication is exposed on the critical path. The ratio compares communication time with ideally partitioned compute time and separates compute-bound from communication-bound clusters (section C.3.3). The compressor and optimizer analysis must establish how bias or variance is controlled, including whether error feedback is required; convergence tests then check that the chosen configuration reaches the target.

The first test is per synchronization. Compression shrinks only the byte term, so it helps a payload whose per-step chunk sits above n^*, such as a full gradient or a large bucket on a small group, and does little for one below it, whose startup term remains (section 5.5). Even above n^*, encoding must cost less than the transfer time it removes.

Napkin Math 5.7: The payback of compression

Problem: A training job takes 40 ms to synchronize a gradient buffer. A compression scheme shrinks the payload, but after scale factors, metadata, and protocol overhead its assumed effective throughput gain is 8×, and encoding and decoding add 2 ms of accelerator time per synchronization. Does compression shorten the synchronization, and by how much?

Math:

1. **Compressed communication time**: 40 ms divided by the effective factor 8 = 5 ms.
2. **Total new time**: 5 ms (comm) + 2 ms (overhead) = 7 ms.
3. **Effective speedup**: 40 ms divided by 7 ms $\approx$ 5.7×.

Systems insight: Compression is a trade of compute for bandwidth. With compression ratio r_{comp} and encoding cost T_{overhead}, it pays off only when $T_{\text{overhead}} < T_{\text{comm}}(N) \times (1 - 1/r_{\text{comp}})$, that is, when encoding costs less than the communication time it removes. In this case, spending 2 ms to save 35 ms of network time is a clear win. On a high-speed NVLink network where $T_{\text{comm}}(N)$ is only 1 ms, the same compression logic would cost more than it saves and slow down training. Always profile the network before adding compression.

The second test spans the full training run, because a compressor that passes the per-step test can still lose once slower convergence adds iterations. Comparing total wall-clock time to the same quality target on a slow and a fast network shows how a modest step-count penalty can reverse the verdict.

Suppose a training run requires 100,000 steps to reach its quality target without compression, with each step taking 500 ms (300 ms compute and 200 ms communication), for a total of 50,000 s. INT8 compression with error feedback of an FP32 payload reduces communication by 4× (from 200 ms to 50 ms per step), but suppose it raises the required steps by 10 percent, to 110,000. The new total is 110,000 steps at 350 ms each, or 38,500 s, a 23 percent improvement. Compression wins here because the per-step saving of 150 ms outweighs the additional steps.

The same model on a faster network, where communication takes only 30 ms per step, reverses the verdict. Compression cuts that to 7.5 ms, saving 22.5 ms per step, but still adds 10,000 steps at 307.5 ms each. The run now takes 33,825 s instead of 33,000 s, 2.5 percent longer. In this serial compute-plus-communication model, the break-even value of ρ depends on the assumed compression ratio and extra steps, with no separate encoding cost included. A training schedule with overlap must instead compare the exposed time saved against encoding and convergence costs. Scheduling transfers during independent computation offers another way to reduce that exposed time.

Checkpoint 5.4: Gradient compression decisions

Verify your understanding of when and how to apply gradient compression:

- ☐ **Error feedback**: explain why uncorrected 1-bit SGD can fail to converge and why 1-bit Adam introduces warmup and residual correction, comparing how each method handles information loss.
- ☐ **Compression payoff**: determine whether applying 4× gradient compression improves total step time when the backward pass takes 800 ms and AllReduce takes either 200 ms or 2,000 ms.
- ☐ **Error accumulator**: derive the steady-state behavior of the error accumulator e_t when the true gradient is constant at $g = 0.3$ and the compression threshold is 0.5.
- ☐ **Top-k vs. quantization**: choose between Top-k sparsification and quantization for a workload, accounting for compression ratio and sparse-index encoding overhead.

5.10 Communication-Computation Overlap

Even with hierarchical routing and gradient compression, waiting until the entire backward pass finishes before launching AllReduce leaves every GPU idle across the full transfer window before the optimizer can update the weights. **Communication-computation overlap** hides that stall by running transfers concurrently with independent backward-pass computation, attacking the communication time that algorithm selection and compression leave exposed. The time it hides is the overlap term T_{overlap} that the Fleet Law (equation 1.3) subtracts from step time. How much of a transfer can be hidden depends on which parts of that transfer occupy the processor and which occupy only the network.

5.10.1 The LogP model of what overlap can hide

The α-β model estimates transfer duration without specifying how much independent computation can proceed during that transfer. Adding the entire estimate to compute time therefore overstates step time when overlap is possible. The **LogP model** (Culler et al. 1993) separates network transit from processor involvement with four parameters:

- L_{LogGP} **(latency)**: The time for a message to traverse the network. The chapter's effective startup cost α bundles this transit time with the processor overhead defined next.
- o **(overhead)**: The processor time spent initiating or completing a transfer. During this work, the same execution resource cannot compute, making o the nonoverlappable cost. In a distributed training context, dispatch, device-kernel launch, memory registration, and runtime progress can all contribute. The value is an effective parameter of a particular software and hardware path, so it must be measured with the same collective, process group, and topology used by the workload.
- g **(gap)**: The minimum spacing between consecutive sends or consecutive receives at a processor. It limits endpoint message rate but does not describe every source of link contention in a routed fabric. The overlap analysis that follows needs L_{LogGP} and o but not the gap.
- N **(rank count)**: The number of ranks in the communication group. The original model names this processor count with its final letter, P; this chapter writes it N to match its other collective cost formulas.

LogP distinguishes network latency (L_{LogGP}, which can be hidden) from processor overhead (o, which cannot). Independent computation can overlap the network transit interval, while initiation and completion overhead remain exposed on the execution resource that performs them. The per-byte transfer time of the α-β model falls on the network side of this split when the network makes progress without the processor.

Napkin Math 5.8: Hiding communication behind computation

Problem: A training pipeline attempts to overlap gradient AllReduce with the next layer's backward pass. The backward pass takes 500 µs. Represent the complete transfer by a toy budget with network time 100 µs, processor overhead $o = 50$ µs to initiate, and $o = 50$ µs to

receive. Assume both overhead intervals occupy the same execution resource as backward computation, with independent network progress and no contention. Can the communication be hidden under these assumptions?

Math:

1. **Overlappable portion**: Network latency L_{LogGP} = 100 µs (data in flight while GPU computes).
2. **Non-overlappable portion**: $2o$ = 100 µs (GPU busy initiating/receiving).
3. **Compute available**: 500 µs.
4. **Hidden**: All 100 µs of L_{LogGP} can overlap with compute.
5. **Exposed**: The 100 µs of $2o$ overhead cannot overlap.

Result: Effective time is 100 µs + max(500 µs, 100 µs) = 600 µs. The network latency is hidden, but the processor overhead remains exposed. Applied to every layer of a backward pass, the same split determines which gradient transfers a training step can hide.

Systems insight: The α-β model captures the total communication time. The LogP model reveals how much of it can be hidden. When designing pipelined training, optimize for low o (kernel bypass, GPUDirect) rather than high β alone.

The α-β model estimates transfer duration from latency and bandwidth measurements. LogP adds the processor/network distinction needed to estimate exposed delay in pipelined execution. Choosing a communication schedule therefore requires both a transfer budget and the independent computation available to overlap it.

When messages are large and pipelined across the network, the LogGP model (Alexandrov et al. 1995) extends LogP with a gap-per-byte parameter G, modeling the streaming cost of long messages. These models describe multiprocessor communication with endpoint rate and capacity limits; they are not restricted to one flow on one link. Their compact parameters do not explicitly represent every shared path, oversubscribed cut, or routing hotspot. For congestion-sensitive workloads, particularly AllToAll for MoE, measurements on the target placement must test whether the model captures the active bottleneck.

↪ The LogGP long-message cost model and its gap parameters are developed in section D.1.2.

5.10.2 Layer-by-layer overlap

Gradient computation during the backward pass proceeds layer by layer, from the output layer back to the input layer. The gradient for layer ℓ is available as soon as that layer's backward pass completes, even while layers $\ell-1, \ell-2, \ldots$ are still computing. That ordering lets the gradient for layer ℓ start communicating immediately while the GPU computes the gradient for layer $\ell-1$. Over a four-layer backward pass, the transfers for layers 4 through 2 run beneath later backward work, and only the transfer for layer 1, which becomes ready last, is left exposed before the optimizer step (figure 5.10).

In PyTorch's `DistributedDataParallel`, this overlap is implemented through gradient hooks. Each parameter registers a hook that marks its gradient ready. Once all required gradients in a bucket are ready, the reducer can launch an asynchronous AllReduce, preserving the same collective order across ranks. Meanwhile, the backward pass continues computing gradients for earlier layers. A transfer is hidden only if it completes before its result is needed; queued communication can consume part of that window, and the last bucket's transfer remains exposed (section 5.10.4).

The effectiveness of this overlap depends on the relative sizes of computation and communication at each layer. For transformer models, the largest gradient tensors belong to the attention and feed-forward weight matrices, which are also the most computationally expensive layers. That computation cannot hide the same layer's gradient transfer, because the gradient is ready only afterward. It can hide a previously launched transfer; the ordering and cost of the remaining layers determine each gradient's overlap window.

The overlap strategy interacts with algorithm choice from section 5.6 and section 5.7. An isolated-collective crossover need not predict the best training-step schedule, because the dependency schedule determines how much of each candidate's transfer time is exposed. Overlap with the

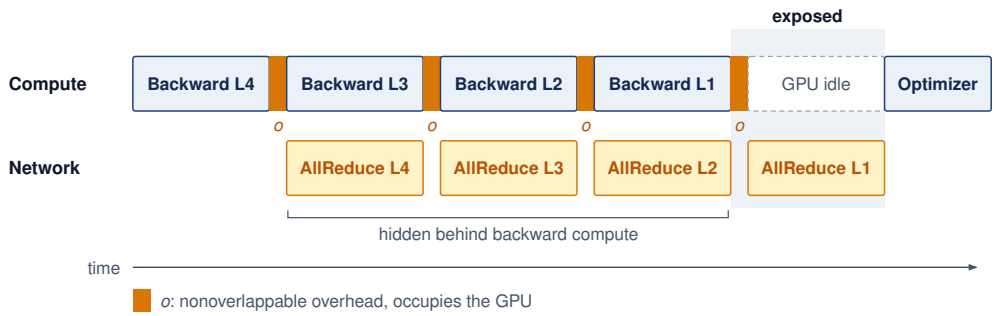

Figure 5.10: Communication-Computation Overlap: The backward pass runs from the last layer to the first, and each layer's gradient AllReduce launches as soon as that layer finishes, so it runs on the network while the GPU computes the next layer's gradients. The first layer finishes last and has no computation left to hide behind, so its AllReduce is exposed before the optimizer step. Overlap succeeds only when the nonoverlappable overhead (o) is small, leaving the bulk transfer free to proceed concurrently. Durations are schematic.

backward pass can reduce Ring's exposed latency penalty, but the benefit must be measured with the actual gradient-ready times and runtime progress behavior. The optimizer step still cannot begin until every required reduction completes, so the last transfer, closest to the input, stays exposed along with any backlog.

5.10.3 Bucket fusion

The layer-by-layer overlap described in section 5.10.2 assumes that each layer's gradient is communicated as a single message. In reality, a transformer layer contains dozens of individual parameter tensors (query, key, value projection matrices, output projection, layer norm parameters, feed-forward weights, and biases). Launching a separate AllReduce for every individual parameter would create thousands of small-message collectives, each paying the full α_{eff} overhead. To avoid this, DDP uses bucket fusion to group parameters into buckets (often configured around tens of megabytes, with 25 MB a common PyTorch reference point) and launches one AllReduce per bucket. The bucket size represents a trade-off: larger buckets amortize α_{eff} overhead but delay the start of communication (because the bucket cannot be sent until all its parameters have computed gradients). Smaller buckets start communication earlier but pay more α_{eff} overhead.

The optimal bucket size depends on gradient-ready times, tensor sizes, and network characteristics. Smaller buckets can start communication sooner when waiting for more gradients would consume the overlap window. Larger buckets can amortize startup when repeated small transfers dominate. Neither choice follows from the model family alone; the useful measure is the resulting training-step time, including the exposed final transfer. PyTorch's `bucket_cap_mb` parameter in DDP allows tuning this trade-off.

Bucket membership also determines how early communication can begin. A bucket containing one late gradient must wait even if its other gradients are ready. DDP can use reverse parameter registration order as an approximation to backward readiness and, depending on configuration, rebuild buckets from the observed backward order. Registration order need not match forward execution order, especially in models with branches or reused parameters. Grouping gradients that become ready together helps preserve the overlap window while maintaining a consistent collective sequence across ranks.

5.10.4 Overlap limits and exposed communication

Even well-tuned bucket fusion cannot hide all communication behind backward compute. Three structural limits (processor launch and completion overhead, boundary stalls at the start and end of backpropagation, and the HBM footprint of in-flight buffers) cap how much transfer time can disappear. The LogP model's overhead parameter o represents processor time consumed by initiating or completing transfers; in the simplified LogP budget, initiation and completion consume $2o$ per

operation on the participating execution resource, so when backward compute intervals are short relative to that overhead, hiding network transit still leaves processor overhead exposed.

The backward pass also has boundary effects. During the first backward computation, no gradient from this pass is ready to communicate. Once the output-side gradient is ready, however, its communication can overlap with the remaining backward work. At the other end, the input-side gradient has no subsequent backward computation to hide its transfer. More layers create a longer overlap window, but depth alone does not guarantee high coverage: if communication takes longer than each compute interval, a backlog remains even in a deep model.

Memory pressure sets the last limit. Buffers used by an in-flight transfer must remain valid while independent computation produces or consumes other tensors. Their overlapping lifetimes can increase peak GPU memory use, although views and buffer reuse can avoid separate copies. For FSDP workloads, prefetching parameters can keep multiple materialized groups resident at once, trading memory savings for communication overlap. The engineering challenge is to hide transfers within the available memory budget, and FSDP's prefetch and resharding policies are the controls that set this trade-off.

Within those limits, per-layer transfer time determines how much of a collective schedule can hide behind backward computation. Pricing a realistic 32-layer transformer backward pass shows how much communication remains exposed when each layer's reduction outlasts its compute window.

Napkin Math 5.9: Overlap budget for a 7B transformer

Problem: A 32-layer transformer model (7B parameters) is trained on 64 GPUs. Assume the data-parallel reducer communicates FP32 gradients, about 875 MB per layer, in 100 MB buckets, four times the common 25 MB default. Each layer's backward pass takes 15 ms. Assume the cluster achieves 3 ms per bucket under its contention and protocol overheads, so each layer's gradients take 27 ms, and assume equal layers and the same charge for the final partial bucket. What is the backward-plus-communication time with and without overlap?

Without overlap (sequential):

$T_{\text{sequential}} = T_{\text{backward}} + T_{\text{comm}}(N)$ = 480 ms + 32 × 27 ms = 1344 ms.

With overlap (pipelined):

Each layer's AllReduce (27 ms) runs in parallel with the next layer's backward pass (15 ms). Since 27 ms > 15 ms, there is 12 ms of exposed communication per layer that cannot be hidden.

Communication runs back to back from the moment the output layer's gradients are ready. Each later reduction overruns its backward window by the exposed amount, so the backlog grows by that much per layer, and the last reduction adds its full duration after the backward pass ends:

$T_{\text{pipelined}} = T_{\text{backward}} + T_{\text{layer comm}} + (N_L - 1) \times T_{\text{exposed}}$ = 879 ms.

Systems insight: Overlap reduces backward-plus-communication time by 34.6 percent. The remaining exposed communication comes from the AllReduce being slower than the per-layer backward pass. To reduce the backlog, either lengthen the independent compute window (for example, with a larger batch) or reduce AllReduce time. The final gradient still has an exposed transfer tail, and any change must preserve the intended training workload and quality target. The overlap is most effective when per-layer backward compute exceeds the reduction time ($T_{\text{layer backward}} > T_{\text{layer comm}}$).

The verdict turns on the assumed per-bucket charge. The hierarchical model of notebook 5.4 prices one uncontended 100 MB bucket on these eight nodes at about 0.85 ms, or 7.6 ms per layer, inside the 15 ms backward window, so at that rate every transfer except the last would hide. An overlap budget is therefore only as good as its achieved per-bucket time, which the calibration of section 5.11.2 measures rather than assumes.

An overlap budget joins the collective's transfer cost to the model's dependency schedule. Realizing that budget requires a runtime whose asynchronous progress and device-memory path preserve the independent work available between gradient production and consumption.

5.11 Communication Libraries and Runtimes

Production libraries implement the ring, tree, and hierarchical schedules, so the engineering decision is which library preserves the chapter's cost, topology, and overlap assumptions on the target hardware.

5.11.1 Backend selection

A collective runtime that bounces GPU tensors through host DRAM over PCIe or ignores NVLink and rail topology can cut effective reduction bandwidth by an order of magnitude even when it implements the right ring or tree schedule. Four properties therefore separate communication backends in practice: the device memory path, topology awareness, portability, and observability. NCCL leads on the first two for NVIDIA GPUs because it executes the collective algorithms of section 5.6 and section 5.7 directly on GPU-resident buffers (Jeaugey 2017; NVIDIA 2026c). Its kernel fusion combines reduction arithmetic with data movement so intermediate results do not bounce back to HBM. GPUDirect RDMA lets the network card read and write GPU memory directly over PCIe without staging payloads through host DRAM. To keep multiple links busy at once, channel pipelining splits each collective across parallel communication channels, though realized throughput still depends on NIC placement, routing, and contention because extra channels cannot create port capacity. Documented environment variables then expose these choices for observability, with `NCCL_ALGO` and `NCCL_PROTO` inspecting or overriding the selected algorithm and protocol and `NCCL_DEBUG` reporting the active topology, channels, and schedule (NVIDIA 2026c).

MPI matters when portability, CPU-side coordination, or HPC integration decides the choice (Message Passing Interface Forum 2015), and two of its features bear directly on overlap. Nonblocking collectives such as `MPI_Iallreduce` return a request handle that the application tests or waits on, a completion model different from NCCL's ordering on CUDA streams, the in-order queues of kernels and copies that a GPU executes. Persistent collectives, introduced in MPI 4.0, initialize an operation once and activate it at every step, amortizing setup for a training loop that repeats the same AllReduce shape. Either feature hides a transfer only when independent work exists and the runtime makes progress. MPI's one-sided communication (`MPI_Put`, `MPI_Get`, `MPI_Accumulate`) maps onto RDMA hardware and suits asynchronous designs in which one process writes another's memory without its participation. CUDA-aware MPI implementations accept GPU pointers, but the selected transport decides whether data still passes through host memory, so that path must be benchmarked against NCCL.

Portability is the strength of Gloo, a PyTorch backend for CPU tensors, development, and testing whose TCP transport exercises collective semantics on one machine before a job reaches an accelerator cluster. Whether a faster shared-memory transport is active must be measured rather than assumed. Other accelerators have their own backends, AMD's RCCL for the ROCm stack and Intel's oneCCL, because the same algorithm has different effective α, β, and topology constraints on each platform. PyTorch's `torch.distributed` routes each call to the backend its process group names, so teams often mix them, with NCCL reducing gradients in a GPU group while a CPU group uses Gloo for barriers, scalar broadcasts, or sampler state. Applied to common scenarios, the four properties select a library, or a pairing of libraries, for each (table 5.13).

Table 5.13: Communication Library Selection: Choose based on hardware and workload requirements. Many NVIDIA GPU training deployments use NCCL; Gloo serves as a portable fallback.

Scenario	Recommended library	Rationale
Multi-GPU training (NVIDIA)	NCCL	GPUDirect, kernel fusion, NVLink-aware
CPU-only distributed training	Gloo or MPI	Mature CPU optimizations
Development/debugging	Gloo	Cross-platform, no CUDA dependency
Mixed vendor GPUs	Vendor-specific groups; validate cross-group path	NCCL, RCCL, and oneCCL are vendor-specific
HPC integration	MPI + NCCL	MPI for job launch, NCCL for GPU collectives

The table follows from the chapter's cost model. A backend's memory path and topology choices determine how it uses the fabric, and its diagnostic signals help identify whether α, β, or processor overhead limits the job. The resulting performance must still be measured with the workload's collective and placement.

5.11.2 Calibrating collective communication models

When an engineer runs the `nccl-tests` benchmark suite on a new cluster and compares the reported `AllReduce` time against the naive point-to-point expression $\alpha + M/\beta$, the achieved bandwidth appears to be roughly half of the link specification even on a healthy fabric. The naive expression omits the collective's traffic factor, about 2 for a ring. Bus bandwidth, defined in section 3.4.4, multiplies the measured rate of tensor size over completion time by the ring's factor whichever schedule ran, so it is a normalization that matches the traffic only for a ring. The bandwidth-latency trade-off (principle 5) provides accurate first-order predictions only when it models the same collective that the benchmark executes. For a ring AllReduce over N ranks and an M-byte tensor, the first-order cost is $2(N-1)\alpha + 2(N-1)M/(N\beta)$, whereas a tree, a hierarchical schedule, and a switch-assisted reduction each carry their own step counts and traffic factors. NCCL also selects protocols, channels, and algorithms according to message size and detected topology, making those runtime choices part of the experimental condition (Jeaugey 2017; NVIDIA 2026c).

A calibrated collective model therefore records the operation, reduction datatype, rank count, tensor size, algorithm and protocol, topology, rank placement, and achieved bus bandwidth. It then fits effective α and β to measurements from that same configuration, as section 3.4.4 does for one measured NCCL AllReduce. Comparing predicted and observed time under this contract reveals software overhead or topology loss without mistaking a collective's algorithmic traffic factor for library inefficiency. The resulting parameters apply only to the measured regime; changing the collective, process group, topology, or library version requires recalibration. A calibrated model selects a collective schedule and sizes its overlap window. Attributing an observed training-step regression to kernels, memory, communication, or contention is a broader measurement task, which the program goodput analysis of section 9.6 develops.

Communication diagnosis begins by comparing the workload against a matched bare-collective baseline and recording the selected algorithm, protocol, rank placement, and achieved bandwidth. A healthy baseline with a slow training step points to invocation granularity, contention, stragglers, or insufficient overlap rather than the fabric alone. Profiler traces and the attribution workflow separate those causes.

The backend must satisfy both the collective contract and the measured execution budget. A fast isolated collective is insufficient when its placement, invocation order, or progress behavior leaves the training step waiting.

5.12 Fallacies and Pitfalls

Communication optimization attracts persistent misconceptions because the interaction between latency, bandwidth, topology, and compression creates nonobvious failure modes that simple mental models miss. Each of the following pitfalls traces a common rule of thumb to the physical or algorithmic boundary where it breaks down.

Fallacy: *Bandwidth is the only metric that matters.*

For small control messages or sufficiently small token exchanges, latency (α) dominates, and buying 400G networking will not help a message that spends 5 µs in a software path slower than the tabulated NDR one (table 5.4). Section 3.4.4 draws that point-to-point split at the critical message size $n^* = \alpha \cdot \beta$. A collective hides a less obvious case. The regime belongs to each schedule step's chunk rather than to the tensor, so a gradient bucket far above n^* can still travel in latency-bound steps on a large ring, where fewer rounds or larger buckets help and added bandwidth helps little.

Pitfall: *Assuming AllReduce works for everything.*

AllReduce computes one reduction over compatible element counts and datatypes; it does not implement destination-specific routing. In expert parallelism and recommendation systems, each worker instead sends distinct data to different peers. These workloads require AllToAll or its

variable-count form, whose performance depends on injection bandwidth, bisection bandwidth, oversubscription, message granularity, routing, and skew rather than logical connection count alone.

Fallacy: *Pipeline parallelism eliminates communication overhead.*

Pipeline parallelism exchanges activations and their backward gradients between adjacent stages. It can coexist with data-parallel AllReduce, and transfer cost still depends on tensor size and placement. Each stage boundary is a point-to-point dependency on the critical path, so a stage cannot start a microbatch until its activations arrive. The pipeline bubble, the time stages sit idle while the pipeline fills and drains, adds to whatever transfer time remains exposed. The bubble analysis of section 6.5.2 shows that shrinking it requires many more microbatches than stages, a constraint on batch size and activation memory that no faster link removes. Parallelism choices must therefore account for both communication dependencies and schedule-induced idleness.

Pitfall: *Using Ring AllReduce by default without checking message size and topology.*

Ring achieves the bandwidth-optimal term $2\frac{N-1}{N}\frac{M}{\beta}$ in the endpoint model but pays $\mathcal{O}(N)$ latency. For a 1 MB gradient across 64 GPUs with $\alpha = 10$ µs and $\beta = 10$ GB/s, the simple Tree model wins because Ring's extra startup cost (about 1140 µs beyond Tree's) exceeds Tree's extra byte-transfer cost (about 1000 µs beyond Ring's). The log-aware crossover estimate $M_{\text{crossover}} \approx N\alpha\beta/\log_2 N$ guides this two-model comparison. Runtime selection should be evaluated with the actual algorithm variants and placement.

Fallacy: *Gradient compression is safe without error feedback.*

Top-k sparsification can discard 99 percent of each transmitted update, but doing so naively can prevent convergence. Without error feedback ($\mathbf{e}_{t+1} = (\mathbf{g}_t + \mathbf{e}_t) - \mathbf{v}_t$), gradients below the threshold are lost forever, accumulating systematic bias that can destabilize training or leave a persistent optimization error.

Pitfall: *Evaluating gradient compression by communication speedup alone.*

Even a 100× communication speedup must be evaluated against the complete training budget when reaching the target quality requires 2× more iterations. It helps only if the reduction in total step time outweighs the extra iterations. Communication microbenchmarks measure collective time, while the engineering decision depends on time to target quality, and the two diverge whenever compression introduces gradient bias or variance that slows convergence. Validate compression on the target convergence benchmark with the deployment training schedule, not on a synthetic AllReduce microbenchmark alone.

Fallacy: *Async collectives always hide latency.*

Python's `dist.all_reduce(..., async_op=True)` only returns control to the CPU. The LogP model distinguishes network latency L_{LogGP} (overlappable) from processor overhead o (nonoverlappable). The condition $T_{\text{compute}} > o$ alone does not guarantee that a transfer is hidden. The remaining independent work must cover the transfer interval, the runtime must make progress, and dependent consumers must wait for completion.

Pitfall: *Assuming network corruption always surfaces as an error.*

Networks are not perfect. A bad cable, faulty NIC, or firmware bug can corrupt data below the level where the training framework sees an explicit error. At 10,000 nodes running 24/7, even rare bit errors can appear often enough to matter. The communication-specific lesson is that bandwidth tuning is incomplete without end-to-end validation: checksums, collective result checks, gradient-norm sanity tests, and the fault-tolerance mechanisms of chapter 7 must catch silent data corruption before corrupted gradients become model updates.

Fallacy: *Flat Ring AllReduce is good enough for multi-node training.*

A single flat Ring AllReduce runs at the rate of its slowest edge. On a multi-node cluster, that edge is one cross-node port at 50 GB/s over InfiniBand rather than the 450 GB/s per direction available within the node, however carefully the ranks are ordered. For an 8-node cluster with 8 GPUs per node, the worked hierarchical schedule assigns one-eighth of the payload to each inter-node shard group and runs the eight groups on eight ports at once. Each port's payload falls 8×, but

the aggregate cross-node bytes do not. The modeled gain of about 4.8× in section 5.7.1 therefore depends on independent port capacity. A multichannel flat schedule that also occupies several ports at once requires its own traffic accounting.

Pitfall: *Assuming one cost formula fits every intra-node algorithm.*

NVSwitch provides high-radix, nonblocking connectivity within supported DGX generations, but it does not turn an arbitrary AllReduce into a one-step reduction. Unless the switch generation and software path provide in-network reduction, endpoints still execute a ring, tree, or other scheduled algorithm over NVLink paths. Newer switch-assisted designs require their own documented cost model. Capacity planning should therefore use the algorithm and hardware generation that the runtime actually selects. Apply the corresponding intra-node traffic factor, the inter-node ring or tree factor, and their hierarchical composition rather than substituting a generic crossbar hop for reduction work.

Fallacy: *Rank-to-GPU mapping does not affect collective performance.*

If the job scheduler assigns ranks to GPUs without considering the physical topology, hierarchical AllReduce and rail-optimized routing may route traffic suboptimally. A common symptom is that `nccl-tests` reaches full bandwidth on the cluster while the training job achieves only a fraction of it, because the job's ranks were placed so that corresponding shards cannot stay on one rail. Message sizes, overlap, and competing traffic can produce a similar gap. Before benchmarking, record each rank's GPU, NIC binding, node, and rail and confirm the intended layout (for example, ranks 0–7 on node 0 and ranks 8–15 on node 1, with local rank i on each node bound to rail i). Contiguous rank numbers on a node alone do not establish that corresponding shards use the intended physical rails. A benchmark that omits this record and the NCCL topology configuration is not reusable evidence, because a placement difference then reads as a library or hardware regression.

5.13 Summary

Every mechanism in this chapter works on the communication term $T_{\text{comm}}(N)$ of the Fleet Law (equation 1.3), either shrinking it through better schedules, placement, and smaller payloads or enlarging T_{overlap} so that less of it is exposed. Each acts on the Communication axis of C^3 so that added compute does not sit waiting. The mechanisms compose into one ordered judgment that starts with the collective contract. The primitive whose output the next computation consumes fixes which ranks exchange data and what each holds afterward, and no faster schedule can recover bytes spent delivering a broader result than the model needs. Once the contract is fixed, the α-β model, the critical size $n^* = \alpha \cdot \beta$, and the ring-tree crossover $M_{\text{crossover}}$ price each candidate schedule round by round on its per-step chunk rather than on the whole tensor.

A schedule priced on a flat network realizes that cost only when hierarchy, rail alignment, and rank order map the logical exchange onto the physical links. Compression and overlap then determine how much of the remaining transfer reaches the critical path. LogP separates the network transit time that overlap can hide from the processor overhead it cannot, and every candidate is judged by time to the same quality target rather than isolated collective speed. A model calibrated against the same collective, process group, and topology closes the loop by separating a schedule's inherent traffic from library or placement loss.

Key Takeaways: Every byte has a travel cost

- **Contract first, then price per round**: Choose the narrowest primitive whose output the next computation consumes, then price its schedule round by round. Each round pays the startup cost on a chunk, so the critical size $n^* = \alpha \cdot \beta$ (principle 5) classifies the chunk rather than the tensor. Collective duration also depends on round count and traffic volume; fit parameters with a benchmark that matches the workload's collective and topology.
- **Algorithm choice is scale-dependent**: Ring AllReduce is bandwidth-optimal but pays $\mathcal{O}(N)$ latency; the simple full-message Tree model uses fewer rounds ($\mathcal{O}(\log N)$) but

more bandwidth work. Use the log-aware crossover estimate $M_{\text{crossover}} \approx N\alpha\beta/\log_2 N$ to reason about the switch point.

- **Hierarchy exposes parallel network capacity**: Local reductions over fast NVLink leave smaller shards for the slow InfiniBand tier. The effective bandwidth gain requires independent port capacity for concurrent shard exchanges; splitting a payload cannot multiply the capacity of a shared uplink.
- **AllToAll schedule selection follows the active bound**: Direct exchange preserves simplicity for balanced coarse chunks, hierarchy amortizes startup across scarce inter-node links, and variable-size exchange removes padding under skew. Injection bandwidth, bisection capacity, oversubscription, message granularity, and the hottest receiver determine which schedule wins.
- **Error feedback controls lossy compression error**: Sparsification can discard 99 percent of each update while accumulating the residual locally and adding it before the next compression step. Under the analyzed assumptions, this bounds accumulated error and restores convergence guarantees even though individual updates can remain biased, but compression pays only if it shortens time to the same quality target.
- **Overlap is limited by remaining independent work**: Gradient hooks and bucket fusion let communication run during later backward computations, but the final transfer and any communication backlog remain exposed. Evaluate the resulting step time, not model depth alone (principle 10).
- **Placement determines usable bandwidth**: Logical collective schedules share physical links according to rank placement and routing (principle 6). A matched benchmark must record that placement; changing the mapping can change the bottleneck even when the algorithm and hardware stay fixed.

If parallelism decides what must be communicated, collective communication measures what that communication costs. Every collective, whatever its contract, is priced first by the same two terms, a startup latency paid once per round and a transfer time paid per byte. The useful optimization is the one that shortens the training dependency path while preserving the required result.

What's Next: From traffic to partitioning

The fleet now has its traffic patterns and a price for each of them: a startup latency per round, a transfer time per byte, and a placement that decides which links carry the load. Which exchanges a training step actually issues, how large they are, and how often they run all follow from how the model and its data are divided among accelerators. The next chapter makes that division explicit. It partitions training across data, tensor, pipeline, and expert dimensions and uses the costs derived here to judge when each partition repays the communication it creates.

↪ Data, tensor, pipeline, and expert parallelism, and the hybrid configurations that combine them, are developed in chapter 6.

6

Distributed Training

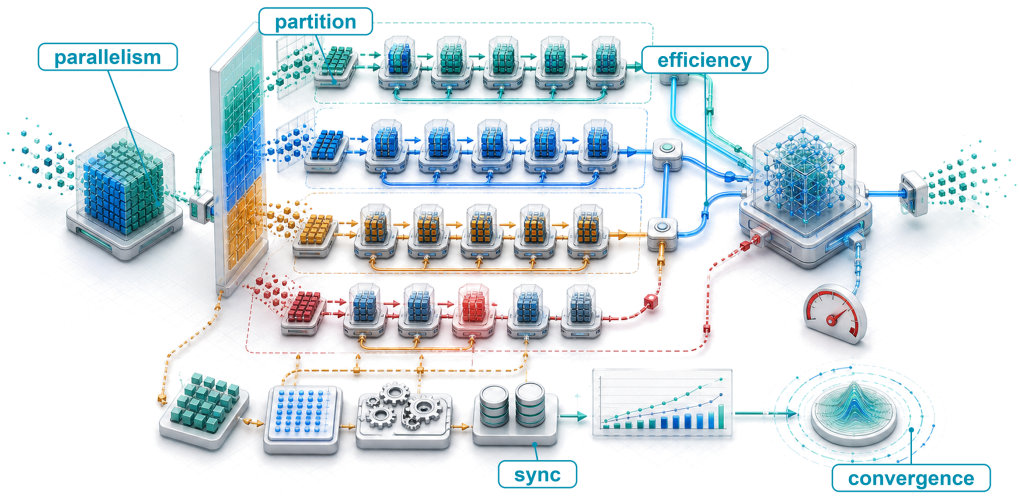

6.1 Why Distribution Is Necessary
6.2 The Parallelism Design Space
6.3 Data Parallelism
6.4 Scaling Efficiency and Convergence
6.5 Model Parallelism
6.6 Hybrid Parallelism
6.7 Multi-Model Training for Alignment and RLHF
6.8 Parallelism Strategy Comparison
6.9 Frameworks and Runtime Contracts
6.10 Fallacies and Pitfalls
6.11 Summary

Purpose

Why can adding accelerators make training slower?

Splitting a training job across machines creates more compute capacity, but it also requires those machines to exchange state and agree on updates. Each accelerator may finish its local work sooner while spending more of the step waiting for other workers. Once communication and waiting consume the time saved by parallel execution, adding accelerators raises the bill without shortening the run. Single-node optimization remains necessary, yet faster local arithmetic cannot remove a dependency on data that another machine has not delivered. Memory pressure complicates the decision, because a model may require distribution simply to fit, even when the resulting exchanges are expensive. Dividing its state across more devices can relieve that limit while exposing another in network bandwidth or scheduling. Conversely, replicating a model can process more examples at once without reducing the memory needed for each copy. Neither improvement alone establishes that training will finish sooner. Changing how examples are grouped or when updates become available can change the work required to reach the target quality. Engineers must therefore compare complete training configurations, accounting for memory, communication, update behavior, and elapsed time together. The practical question is where additional hardware buys useful progress and where it merely creates more coordination. That judgment determines which models and experiments fit the available schedule and budget. In C^3 terms, each partitioning strategy trades parallelized compute against network-bound communication and coordination across workers.

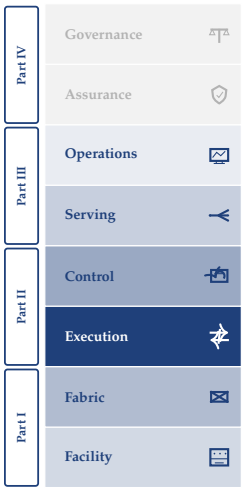

Learning Objectives

- Apply the Fleet Law to diagnose compute-, memory-, communication-, or coordination-bound distributed training regimes
- Calculate scaling efficiency, estimate useful batch size, and budget synchronization overhead for data-parallel training
- Design memory-sharding plans that trade accelerator capacity for additional communication
- Map tensor, pipeline, sequence, expert, and data parallelism onto hardware bandwidth tiers and model structure
- Analyze microbatch pipeline schedules for bubble overhead, activation residency, and synchronous update ordering
- Evaluate synchronization and low-precision policies under SGD equivalence, staleness, straggler, and numerical-stability constraints
- Synthesize hybrid parallelism configurations for frontier, recommendation, and alignment training constraints

6.1 Why Distribution Is Necessary

A GPT-3-class 175B-parameter reference model cannot load onto a single accelerator, because its FP16 weights alone exceed the memory of the device that must hold them. The empirical scaling law (principle 9) keeps pushing frontier models further past what one machine can store or finish in calendar time. The accelerator hierarchy, network fabric, and storage pipeline (chapter 4) form the physical foundation of the fleet, and chapter 5 priced the collective operations that move state across it. The remaining challenge is deciding how to partition the training job: which intermediate results and gradients must cross that fabric, and when workers must synchronize, while preserving the optimization process a single device would have followed.

↩ The power-law relationship between quality and compute, data, and parameters, and the compute-optimal allocation it implies, are established in section 1.3.

In the fleet stack (figure 1.4), distributed training occupies the Distributed Execution row of the Distribution band, the logic that partitions the *mathematical* workload across the *physical* fleet. Data, tensor, pipeline, and hybrid parallelism create different traffic patterns for the interconnect to carry, and selecting among them requires evaluating the reference workload against finite high-bandwidth memory (HBM) capacity, compute throughput, and input bandwidth.

Three distinct limits push a training workload past a single accelerator. Memory exhaustion hits first when model parameters, optimizer states, and activation storage exceed single-device capacity. For full mixed-precision training with Adaptive Moment Estimation (Adam), the convention of 16 bytes per parameter (FP16 weights and gradients, an FP32 master weight, and two FP32 Adam moments) can exceed an 80 GB accelerator at roughly 5 billion parameters before activations. Larger 10–20B models therefore need sharding, offload to a slower memory tier, or another memory-saving approach (Rajbhandari et al. 2020). The capacity limit alone does not choose among them.

↪ Accelerator HBM capacity, the ceiling these per-parameter budgets are measured against, is developed in section A.4.

Even when a model fits in memory, unacceptable training duration can still force distribution because a single device may require weeks or months to converge, turning wall-clock time itself into a systems constraint. GPT-3's 175B-parameter training run used a cluster of V100 GPUs (Brown et al. 2020), illustrating why the calendar, alongside HBM capacity, drives scale-out.

Input demand at dataset scale adds a third pressure. Dataset bytes govern storage capacity, whereas the training schedule governs the required delivery rate. As section 4.4 showed, either can motivate distributed data services, although a large dataset alone does not require distributing model computation. Together, capacity, calendar, and input delivery define the constraints a training configuration must satisfy.

The first signal, memory exhaustion, has a long history, traced in figure 6.1. Over the plotted span, single-accelerator memory grew about 15× from the K80 to the B200, while the memory needed to train a state-of-the-art model grew by roughly 4 orders of magnitude. Among the plotted models, the curves cross in 2020 at GPT-3, whose training state exceeded the largest accelerator then available by about 35×. GPT-3 marks where the gap passed an order of magnitude, and every later plotted

frontier model also sits far above the capacity curve. For those models, spreading state beyond one device's memory is a precondition for training at all rather than a way to train faster.

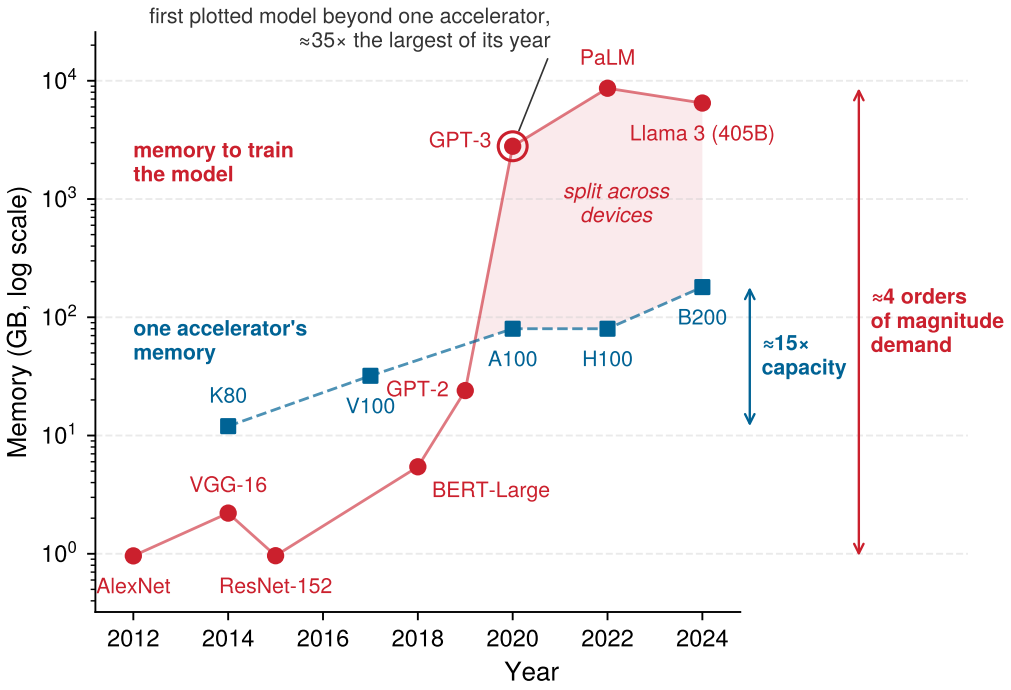

Figure 6.1: The Memory Wall: Accelerator memory capacity plotted against the memory a state-of-the-art model needs to train. Capacity climbs steadily while demand climbs faster, and the widening gap is why weights, gradients, and optimizer state must be split across devices rather than merely compressed. Training memory assumes 16 bytes per parameter for mixed-precision Adam state; parameter counts are from Epoch AI (2026c) and capacities are vendor specifications.

Whichever signal motivates distribution, the system must state which optimization update it implements. The synchronous baseline is developed first. In it, the 1,024 GPUs of the volume's Archetype A (GPT-4/Llama-3) reference cluster combine their contributions before advancing to the next parameter version. The archetype names a class of workload, and the running instance of it in this chapter is a GPT-3-class 175B-parameter dense transformer. Matching the global batch, loss normalization, and optimizer state preserves the corresponding single-device update mathematically. Later synchronization models deliberately relax that update ordering and must be evaluated under their own convergence assumptions.

↩ The three systems archetypes and their limiting constraints are established in section 1.6.1.

Definition 6.1: Distributed training

Distributed training partitions the optimization loop across multiple compute nodes, distributing data or model computation and coordinating contributions according to a specified update protocol.

1. **Significance**: Partitioning training state across devices can meet a memory budget that exceeds one accelerator. GPT-3 (175B parameters) requires approximately 350 GB in BF16 (FP16 also stores 2 bytes per value, so the chapter's byte counts hold for either format), more than 4× the 80 GB capacity of a single H100. If weights remain entirely in HBM, their aggregate capacity lower bound is 5 H100s. Adding BF16 gradients plus FP32 Adam master weights and moments gives about 2.8 TB, or at least 35 H100s before activations. These lower bounds do not establish a runnable placement; offload changes the capacity calculation, and the time budget can require more accelerators than memory alone.

2. **Distinction**: CPU threads can share a coherent address space, while accelerators generally hold separate training-state allocations even within one node. Distributed training extends coordination across those devices and node boundaries through collectives or point-to-point transfers governed by the chosen update protocol.
3. **Common pitfall**: A frequent misconception is that distributed training scales linearly with node count. In practice, communication overhead grows with cluster size and the serial fraction of each step (Amdahl's Law): a fixed serial fraction of the single-device baseline bounds strong-scaling speedup.

Distribution expands three costs that determine whether a partition is viable. Communication overhead includes gradient aggregation through AllReduce, which returns the same reduced tensor to each participant, or rank. For P parameters across N devices, a ring AllReduce, in which each worker passes gradient chunks only to its neighbor in a logical ring, sends approximately $2(N-1)/N \times M_{\text{grad}}$ bytes per worker per step. Here the gradient payload M_{grad} is the P gradient values times their storage width. The ring volume counts one worker's traffic rather than total fleet traffic.

↩ The ring schedule and the derivation of its $2(N-1)/N$ volume are established in chapter 5.

Fault tolerance becomes more demanding as the number of failure sources grows. Under independent constant-rate failures, 1,024 GPUs that each have a mean time to failure (MTTF) of 50,000 hours produce a GPU failure about every 49 hours, the cluster's mean time between failures (MTBF). Whether each event interrupts training depends on recovery and redundancy; section E.1.3 develops the assumptions.

Algorithmic stability adds a separate constraint when the partition changes global batch size or update ordering. Learning-rate scaling and warmup are candidate responses to large-batch behavior (Goyal et al. 2017), not requirements unique to multiple machines. A viable design must keep all three costs within budget.

Communication and coordination costs enter every step through the Fleet Law (principle 10), which decomposes the per-step time:

$$T_{\text{step}}(N) = \frac{T_{\text{compute}}}{N} + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}}$$

Each step takes the single-device compute time divided among the devices, plus the time spent moving data, plus the time spent waiting on coordination such as barriers and stragglers, minus whatever part of that time stays hidden behind computation.

↩ The Fleet Law and its compute, communication, and coordination terms are established in section 1.5.4.

The **communication-computation ratio**, $\rho = T_{\text{comm}}(N)/(T_{\text{compute}}/N)$, compares communication time with ideally partitioned compute time for fixed global work. It is the time consequence of communication intensity CI (equation 1.1), the bytes moved per operation of the workload. When transfers run at link bandwidth β and each device computes at sustained rate $R_{\text{peak}}\eta_{\text{hw}}$, $\rho \approx \text{CI} \times R_{\text{peak}}\eta_{\text{hw}}/\beta$, so the same workload becomes more communication-bound on faster accelerators or slower links. This ratio identifies the pressure to overlap communication; the complete step time also depends on how much exchange is hidden and how much coordination remains exposed.

↩ The α-β message model, with startup latency α and sustained bandwidth β, is established in section 3.4.4.

↪ The communication-computation ratio, evaluated for representative data-parallel and point-to-point transfers and related to communication intensity, is developed in section C.3.3.

Two regimes follow from this ratio. When $T_{\text{compute}}/N \gg T_{\text{comm}}(N)$, the cluster is compute-bound, and communication is small relative to useful computation, provided coordination and input delivery do not introduce another limit. Large-batch training of a dense convolutional model such as ResNet commonly sits here. When $\rho \gtrsim 1$, $T_{\text{comm}}(N)$ matches or exceeds T_{compute}/N and the cluster is communication-bound. Insufficient overlap then leaves substantial time waiting for gradients or activations. Large language models (LLMs) with billions of parameters to synchronize frequently land in this regime. Deep Learning Recommendation Model (DLRM)-style workloads also load the network through their embedding exchanges, but memory capacity and lookup rate often limit them first (section 6.5.6). These are typical operating points rather than fixed labels. Batch size, partitioning, and placement can move the same workload across the boundary, so the ratio must be measured for the configuration at hand. Where a workload lands also depends on how its training loop is cut across devices, the choice each parallelism strategy makes.

↪ The optimizations that help, and the ones that are wasted, for a compute-bound, communication-bound, or coordination-bound fleet are collected for reference in section B.5.2.

6.2 The Parallelism Design Space

Once a model or its training time outgrows a single GPU, the training loop offers three nested axes to split across devices: the examples in a batch, the layers of the model, and the matrix operations inside each layer. Data parallelism assigns different examples to model replicas; pipeline parallelism assigns successive layer groups to different devices; tensor parallelism splits the operations inside a layer. Because every split adds inter-device communication and synchronization overhead on top of single-machine compute and memory limits (Appendix A), profiling and design must treat these three strategies as *loop transformations*.

In compiler terms, data parallelism acts like a parallel for-loop that assigns different batch indices to devices running the same model. Tensor parallelism resembles vectorization, or single instruction, multiple data (SIMD), by partitioning the work inside matrix operations. This is an analogy for the work split rather than for a shared instruction stream or register file, since the devices still exchange tensor shards over an interconnect. Pipeline parallelism instead resembles instruction pipelining, assigning successive layer groups to stages so that overlapping microbatches, the slices of a batch that flow through the stages one after another, keep those stages active, subject to dependency delays and load balance.

However the loop is split, distributed training[1] spreads the workload across machines that must coordinate to train a single model. Coordination here means keeping every model shard or replica on a compatible training step. Basic barriers can keep a small research run ordered, but long-running training jobs also need timeout, checkpoint, and recovery mechanisms so one failed worker does not waste days of compute.

Mapping those three loop transformations onto hardware follows the cluster's interconnect hierarchy. Within a single node, multiple GPUs coordinate across separate device memories over high-bandwidth NVLink or PCIe links,[2] with collective libraries such as NCCL scheduling the transfers.

Crossing node boundaries adds slower inter-node network links, fault-tolerance exposure, and cluster orchestration. Single-GPU kernel and memory efficiency must be tuned before crossing them, because any waste on one device is multiplied across the fleet. Matching the three loop splits to these intra-node and inter-node tiers decomposes the design space onto a three-dimensional grid, $N = N_{\text{DP}} \times N_{\text{PP}} \times N_{\text{TP}}$, that partitions cluster execution across tensor parallelism (N_{TP}, intra-layer sharding), pipeline parallelism (N_{PP}, sequential stage partitioning), and data parallelism (N_{DP}, batch replication). Each cell of the grid is one accelerator coordinate (figure 6.2).

Under the three-dimensional partitioning of figure 6.2, called **3D parallelism**, the total accelerator allocation scales as the product $N = N_{\text{DP}} \times N_{\text{PP}} \times N_{\text{TP}}$. The axes represent distinct mathematical partitions, but they are not independent configuration knobs. Tensor and pipeline degrees must divide the model structure, the data-parallel degree and microbatch schedule must produce a viable global batch, every partition must fit memory, and each communication group must map onto an adequate topology and acceptable failure domain. Training systems therefore select a jointly feasible coordinate rather than tuning N_{DP}, N_{PP}, and N_{TP} in isolation.

6.2.1 Parallelism strategy selection and engineering trade-offs

A parallelism strategy that places frequent, high-volume exchanges on a slow link fails before it leaves the whiteboard, because communication stalls the accelerators it was meant to feed. Choosing a strategy is therefore a constraint satisfaction problem governed by parameter count (P), batch size (B), and interconnect bandwidth. A first feasibility screen compares the traffic each strategy creates against the hardware topology (table 6.1). In that screen, N_L denotes layer count, S sequence length, and d hidden width. The table identifies traffic sources rather than complete byte budgets; datatype width, partition degree, and communication frequency enter the worked mechanisms below.

[1] **DistBelief**: Google's DistBelief (2012) was an early framework for training neural networks across thousands of machines, but its parameter server architecture created bandwidth bottlenecks at central nodes. This limitation drove the shift to decentralized AllReduce patterns in successors such as Horovod and PyTorch DistributedDataParallel (DDP), where each data-parallel worker sends and receives about $2(N-1)/N$ times the gradient size per step rather than concentrating traffic at a single server (Dean et al. 2012; Sergeev and Balso 2018; S. Li et al. 2020; Patarasuk and Yuan 2009).

↪ Checkpointing and recovery for long-running training jobs are developed in chapter 7.

[2] **NVLink**: NVIDIA's point-to-point GPU interconnect delivers 300 GB/s–450 GB/s per-direction bandwidth, roughly 12×–18× InfiniBand HDR per port. This bandwidth gap is why tensor parallelism, which requires AllReduce on every layer, favors the fast NVLink domain, while pipeline and data parallelism can tolerate slower links when their traffic budgets fit. NVLink Switch systems can extend that fast domain beyond one server.

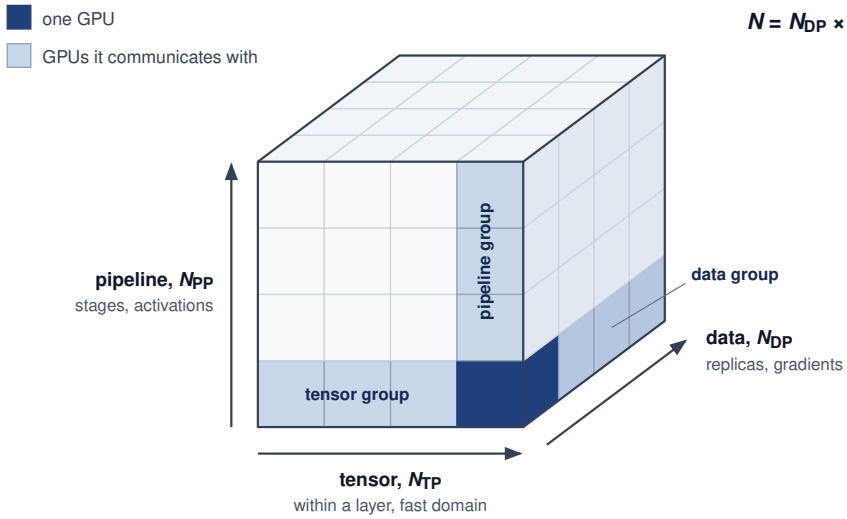

Figure 6.2: The 3D Parallelism Cube: The configuration space decomposes cluster training across three orthogonal axes: tensor parallelism (N_{TP}, intra-layer sharding), pipeline parallelism (N_{PP}, sequential stage partitioning), and data parallelism (N_{DP}, batch replication). Each cell is one accelerator, identified by its $(N_{\text{DP}}, N_{\text{PP}}, N_{\text{TP}})$ coordinate, so the cube holds $N = N_{\text{DP}} \times N_{\text{PP}} \times N_{\text{TP}}$ devices. A feasible point must jointly satisfy model divisibility, per-device memory limits, and the network hierarchy that carries each axis's traffic.

Table 6.1: Parallelism Communication Costs: Tensor parallelism exchanges partial layer results frequently and favors a fast communication domain, which can span servers. Pipeline parallelism transfers boundary activations; its bandwidth need depends on tensor shape and microbatch schedule. Data parallelism synchronizes gradient bytes, whose exposed cost depends on overlap and topology. Memory sharding addresses capacity but does not automatically reduce communication.

Strategy	Communication pattern	Comm. volume	Hardware constraint
Data Parallel (DP)	AllReduce Gradients	$\propto M_{\text{grad}}$ (gradient bytes)	Requires high bisection BW
Tensor Parallel (TP)	AllReduce Activations	$\propto B \times S \times d \times N_L$ (per-layer activations)	Fast domain matched to layer traffic
Pipeline Parallel (PP)	Point-to-Point (P2P)	$\propto B \times S \times d$ (Activations)	Link must sustain stage activations

Those bandwidth requirements impose a hard constraint on hardware placement: each parallelism axis must map onto an interconnect tier that can sustain its traffic within the compute budget. A quick back-of-envelope bandwidth feasibility check catches a mismatch before any code is written.

↩ The communication domains and bandwidth hierarchy used to evaluate placement are established in section 2.5.1.

Systems Perspective 6.1: Bandwidth feasibility check

Tensor parallelism leaves each device with part of a layer's output, and an AllReduce sums those partial outputs twice per transformer block in the forward pass, once after attention and once after the feed-forward network (section 6.5.3). These frequent tensor-parallel collectives must fit the layer-time budget. The test compares the time one layer's local share of computation takes with the time its collectives take on each candidate link. For the forward pass of one GPT-3 layer split 8 ways, processing one 2,048-token sequence at half of an H100's dense FP16 peak:

- Local compute per layer: about 1.9 ms
- Tensor-parallel transfer per layer over NVLink: about 0.4 ms
- The same transfer over one NDR InfiniBand port: about 3.5 ms

On NVLink the exchange is a fraction of the layer's compute; across InfiniBand it would take longer than the compute it serves. The amount hidden depends on independent work available in the execution schedule. NVLink is one suitable fast domain, but the test uses message bytes, exposed latency, and sustained link rate rather than an interconnect brand or chassis boundary. Cross-rack placement may favor pipeline or data parallelism when the available fabric cannot meet the tensor-parallel budget; their activation and gradient traffic must also be checked.

The communication costs in table 6.1 already narrow the choice before any mechanism has been developed, because a strategy is viable only when its traffic fits the tier that must carry it. What each formula implies in practice depends on the mechanics of data, tensor, and pipeline parallelism developed over the rest of the chapter, and section 6.8 turns the same constraints into a decision procedure once those mechanics are in place.

6.3 Data Parallelism

When a model fits on one accelerator but training takes weeks, replicating its computation lets several workers process examples concurrently. The initial data-parallel baseline gives each worker the same model state and a local portion of one global minibatch. Combining those local contributions before the optimizer update preserves the intended global-batch estimator; the sampling and weighting conditions below explain when that equivalence holds.

Definition 6.2: Data parallelism

Data parallelism distributes training examples across workers that execute the model and combine their contributions to its optimization updates.

1. **Significance**: With N workers each processing a local batch of B_{local} examples once per update, the global batch size is $N \times B_{\text{local}}$. At fixed local work, 1,024 workers offer an ideal 1,024$\times$ throughput increase before exposed overhead. A 1B-parameter model with BF16 gradients contributes a 2 GB gradient payload; ring AllReduce moves $2(N-1)/N$ times that payload per worker. Realized speedup depends on local compute, collective service, overlap, and load balance.
2. **Distinction**: The replicated synchronous baseline keeps a full model copy on each worker and aggregates gradients before each optimizer update. Model parallelism instead partitions the model's computation. Sharded data parallelism, developed below, retains data-parallel computation while distributing training-state storage.
3. **Common pitfall**: Data parallelism does not scale indefinitely with worker count. Holding the local batch B_{local} fixed while adding workers grows the global batch. Beyond the workload-dependent critical batch size, additional examples yield diminishing reductions in the steps needed to reach target quality (Shallue et al. 2019). The crossover developed in section 6.4.6 therefore limits the statistical benefit even when gradient variance continues to decrease.

Data parallelism applies when examples can be processed independently, as in training a ResNet model (He et al. 2016) on ImageNet. Each replica computes its local loss gradient, and the workers combine those contributions to reproduce the intended global-batch gradient. The averaging rule makes the sampling and weighting assumptions explicit.

Consider a model with parameters θ training on a dataset D. The loss function for a single data point x_i is $\mathcal{L}(\theta, x_i)$. In standard stochastic gradient descent (SGD) with batch size B, the gradient update for a minibatch is:

$$\mathbf{g} = \frac{1}{B} \sum_{i=1}^{B} \nabla_\theta \mathcal{L}(\theta, x_i)$$

In data parallelism with N devices, each device k computes gradients on its own minibatch $\mathcal{B}_k$:

$$\mathbf{g}_k = \frac{1}{|\mathcal{B}_k|} \sum_{x_i \in \mathcal{B}_k} \nabla_\theta \mathcal{L}(\theta, x_i)$$

When all workers use the same local batch size, the global update averages these local gradients:

$$\mathbf{g}_{\text{global}} = \frac{1}{N} \sum_{k=1}^{N} \mathbf{g}_k$$

Under that equal-batch assumption, the averaging is mathematically equivalent to computing the gradient on the combined batch $B_{\text{total}} = \bigcup_{k=1}^{N} \mathcal{B}_k$:

$$\mathbf{g}_{\text{global}} = \frac{1}{|B_{\text{total}}|} \sum_{x_i \in B_{\text{total}}} \nabla_\theta \mathcal{L}(\theta, x_i)$$

For unequal local batch sizes, the combined-batch gradient is the weighted average $\mathbf{g}_{\text{global}} = (1/|B_{\text{total}}|) \sum_k |\mathcal{B}_k| \mathbf{g}_k$. Data parallelism therefore maintains the statistical properties of SGD training, because sampling the intended global batch across devices, computing local gradient sums independently, and normalizing by the aggregate denominator reproduces its estimator. The reduction step uses an AllReduce over local gradient sums and valid-item counts, the numbers of examples or tokens that contribute to the loss.

Three batch quantities recur from here on, and the chapter keeps them apart. A worker's local batch B_{local} is the set of examples it processes for one update, playing the role of $\mathcal{B}_k$ above. The global batch $B_{\text{global}} = N \times B_{\text{local}}$ is the set the update averages over. Pipeline parallelism later splits a batch into m microbatches of size B_{micro} that flow through stages one after another; the word microbatch is reserved for those pipeline slices.

↪ The ring and tree AllReduce cost models that set the price of this synchronization are developed in section D.2.1.

> **Checkpoint 6.1: Data parallelism mechanics**
>
> Verify your understanding of how data parallelism distributes work:
>
> - ☐ **State placement**: determine whether data parallelism shards or replicates the model state (weights) across GPUs.
> - ☐ **Global batch size**: calculate the effective global batch size for 8 GPUs with a per-GPU batch size of 32.
> - ☐ **Single-device equivalence**: explain why data parallelism is mathematically equivalent to single-device training with a larger batch size.
> - ☐ **Limiting constraint**: identify the primary hardware resource that constrains data parallelism as more nodes are added: GPU TFLOP/s or network bandwidth.

The method parallels gradient accumulation, where a device accumulates contributions from multiple forward and backward passes before one optimizer update. Both techniques use gradient addition and the intended loss denominator to process a larger effective batch. Distribution adds communication, partial failures, and coordination to that arithmetic.

6.3.1 Data parallelism implementation

The SGD equivalence holds only when each phase preserves the intended aggregate sampling and weighting, complete local gradient contributions, and a single synchronized update. The five-stage

execution pipeline in figure 6.3 follows the dataflow across four GPU swimlanes, from assigning local batches to synchronizing the computed gradients. Identical model replicas ingest disjoint batch shards and compute local gradients ($\mathbf{g}_0$–$\mathbf{g}_3$). They then synchronize through a ring AllReduce collective before applying uniform parameter updates.

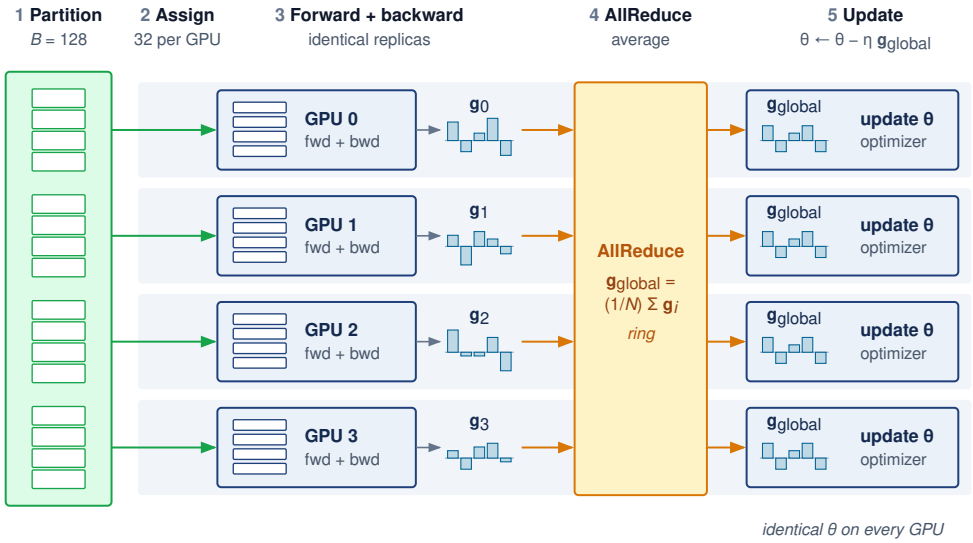

Figure 6.3: Data Parallelism Implementation Pipeline: The five-stage workflow for data-parallel training: (1) partition the global batch, (2) assign sharded mini-batches to individual GPUs, (3) compute independent forward and backward passes to produce local gradients $\mathbf{g}_i$, (4) execute an AllReduce collective across the cluster to obtain the average gradient $\mathbf{g}_{\text{global}}$, and (5) update parameters identically across all devices. This strategy preserves the mathematical semantics of single-device SGD while distributing the compute burden.

The gradient synchronization stage of figure 6.3 (stage 4) creates a dependency, because each rank must receive the completed reduction before applying its optimizer update. The exposed wait can reflect bandwidth, collective startup, or a late contribution. Measuring that wait determines whether synchronization limits scaling; the dependency alone does not identify the limiting resource.

6.3.1.1 Dataset splitting

Data splitting is the first place where the SGD equivalence can fail. The aggregate examples or tokens, their sampling probabilities, and their gradient weights must reproduce the intended global-batch estimator. Unique, equal-sized shards satisfy that invariant, but they are sufficient rather than required. Uniform sampling without replacement can assign disjoint shards. Sampling with replacement can legitimately repeat examples. Variable-length batches can assign unequal token counts if the reduction weights each worker by its valid-example or valid-token denominator, which excludes the padding tokens that fill out shorter sequences. An unweighted mean of unequal local means instead overweights workers with smaller denominators.

The sampler is therefore an integral component of the distributed training system rather than merely an input-loader convenience. A distributed sampler coordinates workers by rank and can deterministically partition or draw indices from a seeded sampling stream. The 1.28 million training images of ImageNet-1k do not divide evenly across 32 GPUs. A sampler that pads by repeating whole examples, unlike the token padding of variable-length batches, gives each GPU 40,037 examples per epoch and adds 17 repeats across the job. The loss normalization or sample weights must account for those repeats when exact epoch-level weighting matters. A recorded seed and deterministic index stream make sample assignment repeatable. Bitwise run reproducibility additionally depends on kernel determinism, collective reduction order, numerical precision, and checkpointed random-number-generator state.

6.3.1.2 Compute-phase forward and backward passes

For a replicated model whose local forward and backward operations require no cross-worker state, the computation phase is *embarrassingly parallel*. Each GPU executes the model on its local batch; synchronized normalization or additional model partitioning can introduce communication within that phase. To isolate the activation term, consider an unpartitioned 175B reference model at a local batch of 4 sequences and sequence length 2048. This is a hypothetical memory budget, not a feasible one-GPU training configuration, because weights and optimizer state must fit as well. Without optimization, storing these activations for backpropagation would consume roughly 1.1 TB of HBM, an order of magnitude beyond the capacity of even an H100 GPU. Activation checkpointing recomputes selected activations during the backward pass rather than retaining them, reducing this example's footprint to ~19.3 GB at the cost of extra computation.

The capacity test must include activations rather than stopping at bytes per parameter. Section A.6 tabulates the 16-byte-per-parameter training state from which this test starts. At the point of peak residency,

$$M_{\text{total}} = M_{\text{weights}} + M_{\text{gradients}} + M_{\text{optimizer}} + M_{\text{activations}},$$

with communication buffers and temporary workspaces consuming additional headroom. Sharding (section 6.3.3) reduces the first three terms according to its stage, whereas sequence length, local batch size, and pipeline scheduling govern $M_{\text{activations}}$. Activation checkpointing reduces that last term by retaining only selected boundaries and recomputing the omitted forward operators during backpropagation (Narayanan et al. 2021). Full recomputation re-executes about one extra forward pass per step. Because a backward pass costs roughly two forward passes, a step that needed three forward-pass equivalents of arithmetic now needs four, about a third more (Chen et al. 2016). Selective recomputation re-executes only operators whose activations are large but cheap to produce, such as the attention softmax and dropout, and keeps most of the memory saving for a few percent of extra arithmetic (Korthikanti et al. 2023). Recomputed operations raise hardware FLOPs utilization (HFU) but not model FLOPs utilization (MFU), which credits only the model's own forward and backward work. A plan that buys feasibility with full recomputation therefore pays for it in useful throughput. The memory saved is purchased with additional FLOPs and a longer critical path. A feasible parallelism plan must budget both effects.

As the backward pass traverses the computation graph in reverse, it produces gradients for the parameters held by each replica or shard. Under pure data parallelism, a full 175B FP16 replica would imply a 350 GB gradient tensor per worker, which exceeds single-accelerator memory budgets. Large runs therefore combine data parallelism with sharding, tensor parallelism, or pipeline parallelism. Each local gradient contributes to the intended global-batch estimator. Aggregation can begin as gradient tensors become ready, overlapping later backward computation, but the optimizer must wait for the completed contributions it consumes. This readiness schedule connects local computation to collective communication.

6.3.1.3 Gradient synchronization

Gradient synchronization combines local contributions into the synchronous update. Each replica must receive the aggregate before applying that update. In a centralized parameter-server design, where servers hold the parameters and workers send them gradients, fixed per-worker payload and update frequency give the server $\mathcal{O}(N)$ incoming traffic. That concentration becomes a bottleneck when demand exceeds the server tier's service capacity; sharding the server state or changing the communication pattern can relieve it.

The replicated synchronous baseline distributes aggregation through collective communication. AllReduce lets each worker contribute its local gradient tensor and receive the same reduced result, which the training loop normalizes to obtain the global average. In sharded variants such as ZeRO and FSDP, ReduceScatter and AllGather change which result each worker retains, leaving each worker only its own shard of the reduced gradient and reconstructing full parameters on demand. Small messages emphasize startup latency; large messages emphasize sustained bandwidth. Centralized aggregation remains another design choice, particularly for the sparse workloads developed later.

↩ The result contracts of AllReduce, ReduceScatter, and AllGather, and the decomposition of AllReduce into a ReduceScatter followed by an AllGather, are established in section 5.4.1.

At the Distributed Execution row of the fleet stack, the loop specifies an AllReduce contract while its implementation determines the cost. For a large worker group, ring AllReduce transfers approximately $2M_{\text{grad}}$ bytes per worker for an M_{grad}-byte gradient payload. Chapter 5 derived

this approximate $2M_{\text{grad}}$ volume and compared algorithms whose startup costs, topology, and contention differ. That chapter also examined how libraries such as NCCL schedule exchanges within and across communication domains.

Because those algorithms predict different times on the same hardware, diagnosing a measured collective starts by comparing its trace against each prediction. A typical comparison sets a flat ring against a hierarchical AllReduce, which reduces within each NVLink node before crossing the inter-node fabric. Consider a hypothetical A100 cluster with HDR InfiniBand running an AllReduce of a 3 GB gradient tensor across 128 nodes (1,024 GPUs). The flat-ring estimate is 239.8 ms, the idealized hierarchical estimate is 48.5 ms, and the trace measures 100 ms. The observation rules out the flat schedule, and the residual above the hierarchical ideal must come from lower achieved rates, startup, or a late participant. Switch counters showing inter-node links at 60 percent utilization do not settle the question, because an aggregate can hide one saturated port during the exposed interval. Per-port and per-channel traces must locate the bottleneck before any topology change.

Across the wider design space, figure 6.4 compares three synchronization topologies. The bandwidth-optimal ring AllReduce circulates chunks along neighbor links ($2(N-1)/N$ volume). The centralized parameter server is constrained by network interface card (NIC) saturation at the server node. The dense mesh of logical AllToAll peer traffic requires $N(N-1)$ messages. In dense synchronous data-parallel settings, ring AllReduce avoids a single reducer bottleneck by distributing traffic evenly across participating links, as Baidu's implementation illustrates (Gibiansky 2017).

↩ Hierarchical AllReduce and its cost relative to a flat ring are established in section 5.7.1.

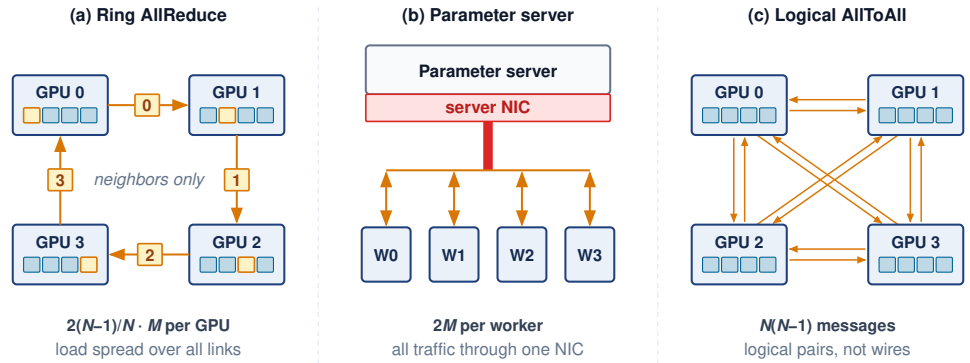

Figure 6.4: Communication Patterns and Their Costs: Ring AllReduce and a parameter server illustrate two ways to aggregate updates. AllToAll instead routes distinct data to each peer; the complete graph shows logical destinations, not a requirement for dedicated physical links. A switched fabric can carry this traffic, with cost determined by message size, routes, and contention.

The panels separate an aggregation algorithm, a centralized architecture, and a collective traffic pattern. They do not provide interchangeable implementations of the same operation. Only the first two aggregate gradients; the AllToAll panel previews the exchange that sparse embedding tables and mixture-of-experts layers need, in which every worker sends different data to every other worker (section 6.5.5). Libraries choose algorithms and routes for the requested collective on the available fabric; a logical complete graph does not guarantee one-round physical delivery.

6.3.1.4 Parameter updating

Parameter updating closes the replicated synchronous contract. After the gradient AllReduce (Patarasuk and Yuan 2009), every worker holds the same global gradient and applies it to its own replica with SGD, momentum, or Adam, so the update stays decentralized without losing equivalence to single-device training. Identical gradients alone are insufficient, because optimizer moments, learning rates, and step counts must also match across replicas. Frameworks automate the cycle of splitting data, computing gradients, synchronizing, and updating, but they cannot remove the ordering constraint that keeps the replicas coherent. Every worker must finish aggregating the current step's gradient before any replica computes from the updated weights.

6.3.1.5 Synchronization models

Whether every worker must wait at a barrier on every step defines the job's synchronization model. The baseline is bulk synchronous parallel (BSP), the superstep-and-barrier model, in which workers compute from one parameter version, combine their gradient contributions, and complete the update before computing from the next version. This is a dependency order, not a requirement that every device update at the same wall-clock instant.

↩ Bulk synchronous parallel execution and its barrier cost are established in section 3.2.

With the same global batch and reduction, BSP matches the corresponding synchronous single-device update mathematically. It does not by itself guarantee bitwise reproducibility. Floating-point reduction order, nondeterministic kernels, data-order state, and failure recovery can still change the result. The synchronization cost is that the slowest worker determines iteration time. A worker that keeps producing correct results but runs markedly slower than its peers is a **straggler**, and under BSP every straggler delays the entire step. Relaxing the barrier changes the optimization process as well as the waiting time.

Stale synchronous parallel (SSP) relaxes this constraint by allowing workers to proceed up to s iterations ahead of the slowest worker before blocking. This bounds staleness while reducing synchronization delays. SSP requires careful learning rate tuning since workers compute gradients on slightly different parameter versions. The bounded staleness guarantee provides a middle ground between BSP's strong consistency and fully asynchronous approaches (Ho et al. 2013).

Asynchronous parallel (ASP) execution removes the global update barrier, allowing a worker's contribution to be applied while others continue computing; asynchronous SGD is the optimizer run under that protocol. Less barrier waiting can improve update throughput, but communication and shared-state contention remain. Gradients computed from older parameter versions introduce staleness that can affect convergence. Section 6.4.4 develops the objective and delay assumptions needed to evaluate that trade-off. The three models lie on one spectrum, from the strong consistency of BSP to the barrier-free updates of ASP (table 6.2).

Table 6.2: Synchronization Model Trade-Offs: BSP, SSP, and ASP trade consistency for throughput along a spectrum. BSP preserves synchronous update boundaries but pays for the slowest worker; SSP permits bounded progress ahead of slower workers; ASP removes the global update barrier without a finite protocol-wide staleness bound. Throughput and time to target quality depend on the workload and implementation. Reproducibility requires a separate deterministic-execution contract.

Model	Consistency	Throughput	Convergence	Use case
BSP	Strong	Bounded by slowest worker	Same synchronous update semantics	Controlled final training runs
SSP	Bounded staleness	Less global-barrier waiting	Depends on delay and tuning	Hyperparameter search
ASP	Weak	No global update barrier	Depends on staleness and objective	Large heterogeneous clusters

Execution traces for four workers with one straggler make the differences visible (figure 6.5). Under BSP, global barriers force the faster workers into idle wait blocks while straggler W3 completes. Under SSP, bounded slack ($s = 1$) lets them run ahead and throttles the drift with brief pauses. Under ASP, asynchronous updates remove the global barrier without imposing a finite staleness bound, although communication and resource contention can still leave workers waiting.

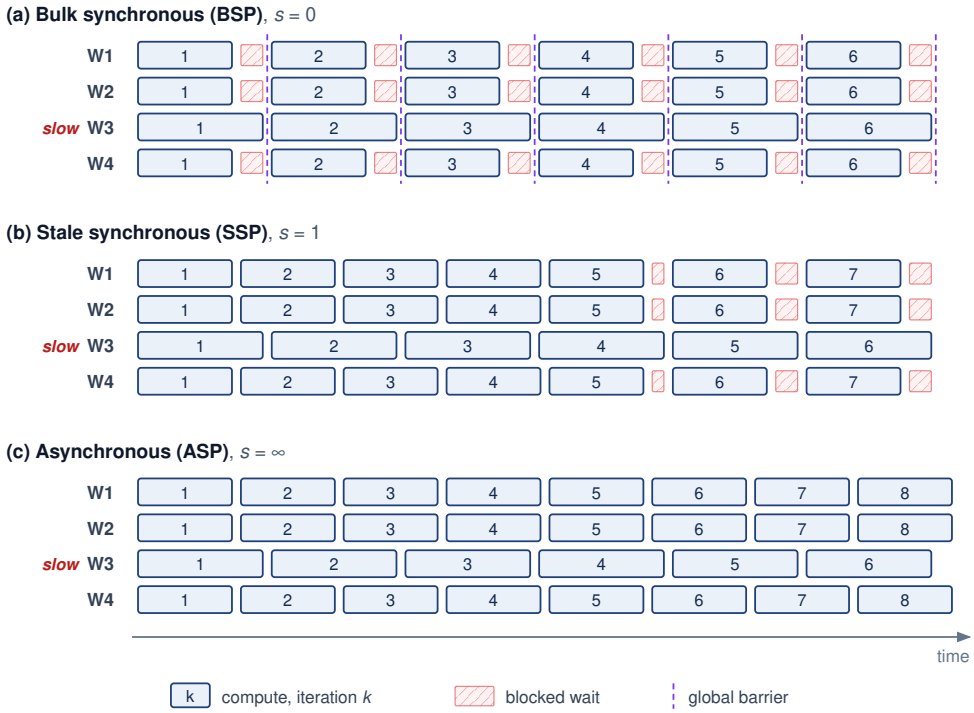

Figure 6.5: Distributed Synchronization Models: The timelines compare three synchronization strategies. Top: bulk synchronous parallel (BSP) forces all workers to wait at a global barrier every step. Middle: stale synchronous parallel (SSP) allows workers to proceed up to s steps ahead of the slowest worker. Bottom: asynchronous parallel (ASP) execution removes the global update barrier while allowing gradient staleness; communication and resource contention can still delay progress.

6.3.1.6 Barrier semantics and failure modes

An optimizer update that consumes an AllReduce result must wait for every required contribution. Workers can execute independent computation while a collective is in flight, but they must issue compatible collectives in the same group order and wait for completion before consuming the result. This coupling creates failure modes absent from single-device training.

A worker failure can leave peers waiting for a contribution that will never arrive. The runtime may report a communication error, or the operation may remain blocked until a timeout detects the failure. Watchdogs bound that wait so the job can enter recovery rather than retain an unusable allocation indefinitely.

Collective mismatches occur when workers issue incompatible operations, tensor shapes, or ordering within a communication group. Conditional computation can create such a mismatch if some workers omit an expected reduction. Variable sequence lengths or different expert routes do not inherently violate the protocol; the runtime must construct a compatible communication schedule for those varying local computations.

Stragglers extend an iteration when their late work lies on its critical path. Thermal throttling, network congestion, or OS jitter can delay a required contribution even when other workers are ready. At 1,000 GPUs with a 1 percent straggler probability per GPU per step, the expected count is

One missing worker stalls the entire AllReduce barrier.

10 stragglers per step, but the number that matters under BSP is the chance of a step with none at all. That chance is the per-GPU on-time probability raised to the GPU count, about 0.004 percent. Almost every step therefore waits for at least one slow worker, which is why the slowest worker, not the typical one, sets step time. The resulting exposed wait depends on how late that contribution arrives and what independent work remains. Deployed systems bound these waits with timeouts and heartbeat monitoring, and elastic training lets a job continue with a resized worker group after a failure.

↪ Heartbeat and collective-timeout failure detection, and the trade-off between restarting healthy workers and leaving peers blocked, are developed in section 7.9.1.

6.3.2 Data parallelism trade-offs

Data parallelism is the simplest way to add accelerators because it leaves each replica's model computation intact, but adding workers eventually hits either a communication wall or a statistical wall (Ben-Nun and Hoefler 2019). Near-linear throughput scaling requires local compute efficiency, balanced input delivery, and small exposed synchronization costs. Which wall arrives first depends on whether the global batch stays fixed or grows as workers are added. Holding the global batch fixed while adding workers is **strong scaling**: each worker's local batch shrinks, so communication must hide behind less computation. Holding the local batch fixed so that the global batch grows with N is **weak scaling**: per-worker computation stays constant, but each update averages more examples, which eventually stops reducing the number of updates needed. The rest of the chapter refers to the two regimes by these names.

Weak scaling pays off when the larger global batch still reaches the target accuracy. Goyal et al. (2017) trained ResNet-50 on ImageNet in about one hour on 256 GPUs at a global batch of 8,192 images; the learning-rate rules that let that batch reach the target accuracy are developed in section 6.4.5.

Replicated data parallelism retains the model's local operators, and frameworks can trigger gradient reductions as backward computation produces them. Each replica executes a complete forward and backward pass, avoiding the bubbles of pipeline parallelism, the idle slots while layer stages fill with work and drain. Input stalls, imbalance, and exposed reductions can still leave devices idle.

These advantages hold only until replication runs into memory capacity, communication bandwidth, or statistical diminishing returns. The memory wall arrives first on large models because a replicated worker holds every parameter, gradient, and optimizer state, demanding more than 1 TB per replica for the 175B-parameter example before activations. Even when the state fits, communication limits useful scaling once its exposed cost grows relative to local work. As N increases, the ideal ring bandwidth term $\frac{2(N-1)}{N}\frac{M_{\text{grad}}}{\beta}$ approaches a constant for a fixed gradient payload M_{grad} and link bandwidth β, while startup latency, topology, and shrinking local computation worsen the balance. Statistical efficiency then sets the ceiling under weak scaling, because growing $B_{\text{global}} = N \times B_{\text{local}}$ with N eventually yields diminishing reductions in update count. Strong scaling avoids that batch-size ceiling only by shrinking each worker's local batch, which exposes the communication trade-off sooner.

Napkin Math 6.1: GPT-2 data parallel scaling: Single node

A GPT-2 scaling scenario makes the efficiency loss concrete by holding the model fixed and changing only the communication domain. GPT-2 stands in for the 175B reference here because it fits on one GPU, so pure replicated data parallelism applies and communication is the only variable. This notebook and the next report parallel efficiency, the achieved speedup divided by the GPU count, which section 6.4.1 develops as scaling efficiency. The first case keeps all eight GPUs inside a single NVLink node.

Single-GPU baseline

- Batch size: 16 (with activation checkpointing, fits in 32 GB)
- Time per step: 1.8 s
- Baseline time for 50K single-GPU steps: 25 hours

The following durations hold the total sample budget fixed, so larger global batches require fewer optimizer updates. They assume those samples attain the same quality; the calculation establishes throughput scaling, not convergence equivalence.

8 GPUs: single node with NVLink

- Per-GPU batch: 16, global batch: 128
- Gradient synchronization: 5.2 GB @ 450 GB/s (NVLink, per direction) $\approx$ 11.7 ms

Performance results:

- Compute: 1800 ms per step
- Communication: 11.7 ms per step
- Total: 1811.7 ms per step
- Speedup (throughput): 8×
- Parallel efficiency: 99.4 percent

Training time: 25 hours ÷ 8 = 3.1 hours

Inside the node, NVLink keeps gradient exchange small relative to compute, so efficiency stays near 99.4 percent. The picture inverts once the same model must synchronize across nodes over a commodity network.

Napkin Math 6.2: GPT-2 data parallel scaling: Commodity scale-out

The second case scales the same GPT-2 run across multiple nodes, replacing the intra-node NVLink hop with inter-node Ethernet for the bulk of the AllReduce. The 10 Gb Ethernet link models a commodity cloud tier without a dedicated training fabric, far slower than the InfiniBand of the reference cluster, which makes the cost of leaving the node easy to see.

Commodity network configuration: 32 GPUs across 4 nodes

- Per-GPU batch: 16, global batch: 512
- Intra-node communication: 11.7 ms (NVLink)
- Inter-node communication: 5.8 GB @ 1.25 GB/s (10GbE) $\approx$ 4650 ms

Performance results:

- Compute: 1800 ms (27.9 percent of time)
- Communication: 4661.7 ms (72.1 percent of time), so communication dominates and becomes the bottleneck.
- Total: 6461.7 ms per step
- Speedup (throughput): 8.9× faster $\rightarrow$ 2.8 hours
- Parallel efficiency: 27.9 percent

The scale-out run bought its larger global batch with four nodes on a slow network. A single node can reach the same global batch another way, by running several local microsteps and accumulating their gradients before one AllReduce. The comparison below is therefore between two placements, one node with accumulation against four nodes without it. Accumulation's job is to keep the global batch equal, while the saving comes from never leaving the NVLink domain.

Napkin Math 6.3: Gradient accumulation speedup

Problem: A GPT-2 run on a commodity 10G network is communication-bound at 32 GPUs across 4 nodes, costing $359 for a fixed number of samples. Can a single 8-GPU node that uses gradient accumulation to match that global batch process the same samples for less?

Math:

1. **Effective batch size**: 8 GPUs $\times$ batch 16 $\times$ 4 accumulation steps = 512.
2. **Communication overhead**: With 4-step accumulation, the job AllReduces once every 4 steps (overhead = 11.7 ms / (4 $\times$ 1800 ms) $\approx 0.162\%$).
3. **Training duration**: Total time is 3.1 hours.
4. **Scale-out baseline cost**: 4 nodes $\times$ 2.8 hours $\times$ \$32/hr = \$359.
5. **Accumulation cost**: 3.1 hours $\times$ \$32/hr = \$100.

Systems insight: The one-node plan saves \$259 (72.1 percent). Almost all of that saving comes from placement, because the four-node run spent most of each step on the slow inter-node link and the single node never uses it. Gradient accumulation is what makes the single node a fair substitute, because it reproduces the four-node global batch. Its own contribution to speed is small, since the same node already ran at 99.4 percent parallel efficiency without accumulation. Accumulation on a multi-node job still requires inter-node synchronization once per window.

In this GPT-2 scenario, changing the communication domain drops efficiency from 99.4 percent to 27.9 percent. Accumulation can preserve the effective batch while avoiding the slow scale-out link, at the cost of additional local computation between updates. It leaves each worker's replicated model state unchanged, however; reducing that memory burden requires sharding.

6.3.3 Memory-efficient data parallelism with ZeRO and FSDP

Replicated data parallelism stores an identical copy of weights, gradients, and optimizer state on every GPU, so a 7B-parameter model demands 112 GB on each device and overflows an 80 GB accelerator before a single activation is allocated. The Zero Redundancy Optimizer (ZeRO) (Rajbhandari et al. 2020) and PyTorch's Fully Sharded Data Parallel (FSDP) (Zhao et al. 2023) remove that duplication by sharding training state across workers while keeping each worker's full forward and backward computation.[3] Sharding can make a model fit in memory without splitting its layer computations across devices.

[3] **ZeRO (Zero Redundancy Optimizer)**: Introduced by Microsoft researchers in 2019 and released in the DeepSpeed library (Rasley et al. 2020), ZeRO's stages progressively partition optimizer states, gradients, and parameters. FSDP draws on the fully sharded approach associated with ZeRO's deepest stage; the shared partitioning principle does not imply identical implementations or communication schedules.

Definition 6.3: Sharded data parallelism

Sharded data parallelism partitions optimizer state, gradients, and at the deepest stage the parameters themselves across data-parallel workers. Each worker processes its own minibatch shard and uses collectives to obtain the parameters and reduced gradients it needs. Full sharding reduces each worker's persistent state toward $1/N$ of the replicated training state; activations and temporary buffers still contribute to peak memory.

1. **Significance**: Under the mixed-precision Adam convention used here, 16 bytes of state per parameter gives a 7B-parameter model 112 GB of training state, exceeding a 80 GB accelerator before activations. ZeRO Stage 3 across 64 workers reduces persistent state per device to 1.75 GB (0.25 bytes per parameter). The resulting AllGather and ReduceScatter traffic trades memory capacity for communication whose exposed cost depends on the interconnect and available overlap.
2. **Distinction**: Unlike vanilla data parallelism, which replicates training state, sharded data parallelism partitions its storage. Unlike model parallelism, it retains each worker's full forward and backward computation. Fully sharded execution gathers parameters for the current computation; the resharding schedule determines when it releases them.
3. **Common pitfall**: A frequent misconception is that sharding is free memory. The on-demand AllGathers place parameter traffic on the critical path of every layer in every step. On slower interconnects or at small per-worker batch sizes, the exposed communication erodes throughput faster than the memory savings help. Capacity is purchased with bandwidth.

The scale of the savings is clearest on a model whose replicated training state only modestly overflows one accelerator, so that the effect of each sharding stage stays visible. The 175B reference returns in section 6.5, when the question becomes whether sharding alone is enough.

Napkin Math 6.4: ZeRO memory savings

Problem: Training a 7B parameter Llama-2 model in mixed precision requires 16 bytes per parameter for weights, gradients, and optimizer state. Does the full training state fit on a single A100-80 GB, and how does ZeRO Stage 3 with 64 GPUs change the per-device memory requirement?

Baseline: Standard DDP (replicated state) Per-Parameter Memory Cost:

- **Weights (FP16)**: 2 bytes
- **Gradients (FP16)**: 2 bytes
- **Optimizer state (FP32)**: 12 bytes (4 master weight + 4 momentum + 4 variance)
- **Total**: 16 bytes/parameter

Total Memory for 7B Model:

$$C_{\text{state,total}} = 7 \times 10^9 \times 16 \text{ bytes} \approx 112 \text{ GB}$$

Baseline outcome: Out-of-memory (OOM) on A100-80 GB.

Optimization: ZeRO-3 (fully sharded) With $N = 64$ GPUs, state is partitioned:

- **Weights**: $2/N$ bytes
- **Gradients**: $2/N$ bytes
- **Optimizer**: $12/N$ bytes
- **Total**: $16/N = 0.25$ bytes/parameter effective storage

Per-GPU Memory:

$$C_{\text{state,ZeRO3}} = \frac{112 \text{ GB}}{64} \approx 1.75 \text{ GB}$$

Result: The persistent state fits, leaving ~78 GB for activations, gathered parameters, temporary buffers, and runtime allocations. Their simultaneous peak determines whether the chosen batch size fits.

ZeRO reaches that result in stages, and the intermediate stages matter because each one buys memory with a different communication pattern. Optimizer sharding alone (ZeRO-1) already brings the replicated footprint of the same model under the accelerator ceiling, because optimizer state is three-quarters of it. Gradient sharding (ZeRO-2) and parameter sharding (ZeRO-3/FSDP) continue toward the fully sharded figure (figure 6.6). Each stage also changes which collectives the update uses (table 6.3).

Table 6.3: ZeRO Sharding Stages: Each stage shards a different layer of training state and trades a different communication primitive. ZeRO-1 and ZeRO-2 both replace the gradient AllReduce with a reduction to each shard's owner and a parameter AllGather, matching DDP volume, and differ only in whether each worker keeps a full-size gradient buffer; ZeRO-3 (FSDP) adds per-layer AllGathers during forward and backward passes, unlocking near-linear memory scaling at $1.5\times$ communication.

Stage	What is sharded	Memory reduction	Communication overhead
ZeRO-1	Optimizer states only	~4×	Reduce to owner + AllGather (1× DDP volume)
ZeRO-2	+ Gradients	~8×	ReduceScatter + AllGather (1× DDP volume)
ZeRO-3/FSDP	+ Parameters	~N (linear in workers)	AllGather before each layer (1.5× DDP volume)

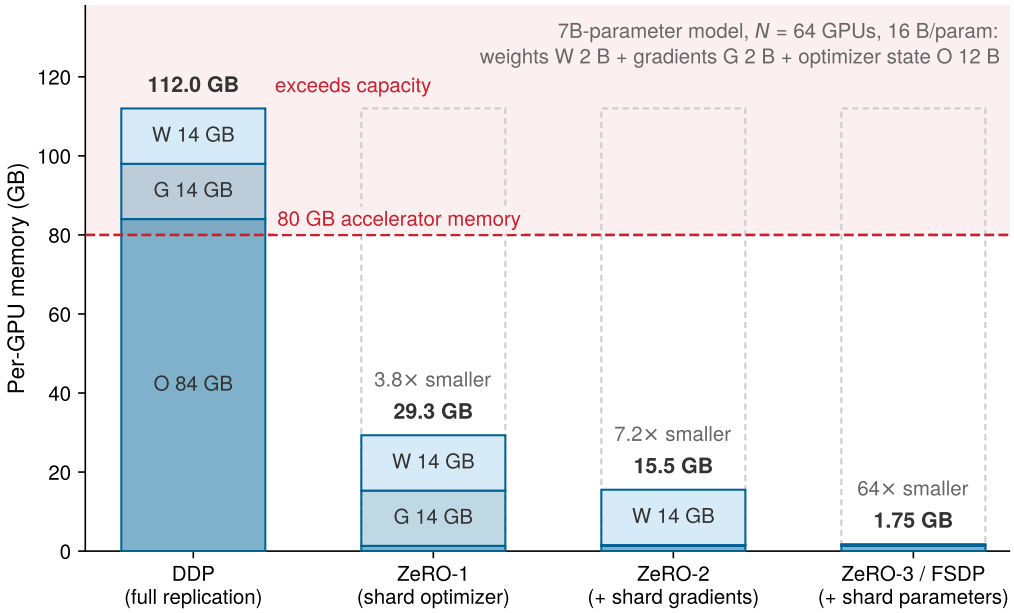

Figure 6.6: ZeRO Memory Reduction: Standard data parallelism (DDP) replicates all model states across every GPU. ZeRO progressively partitions these states: ZeRO-1 shards optimizer states, ZeRO-2 adds gradient sharding, and ZeRO-3 shards the parameters themselves; each dashed outline marks the replicated stack a stage no longer holds. ZeRO-3 achieves linear memory scaling, enabling models with 100B+ parameters to fit on accelerators that could not hold the replicated state.

The first step, ZeRO-1, shards optimizer states across GPUs so each GPU stores only $1/N$ of the Adam optimizer-related state. During backpropagation, each worker still materializes the full 2 bytes/param gradient buffer locally, which is why gradient memory remains unsharded in table 6.3. Once gradients are ready, however, they need only be reduced to the worker responsible for each parameter shard via a ReduceScatter, and after each worker updates its owned shard, an AllGather broadcasts the updated parameters to restore the full replicas. This two-phase schedule (one gradient ReduceScatter plus one parameter AllGather) matches the table's $1\times$ baseline communication volume, whereas running a full gradient AllReduce followed by a parameter AllGather would add 50 percent more traffic. Under the 12-byte convention that counts FP32 master weights, momentum, and variance, sharding the optimizer state reduces cluster-wide storage from $12N$ bytes/param to 12 bytes/param in total, split evenly across the GPUs.

Building on that owner-update contract, ZeRO-2 eliminates persistent full-model gradient buffers by streaming the ReduceScatter during the backward pass. As each layer or communication bucket finishes its backward computation, workers immediately ReduceScatter that bucket's gradients and free the slices owned by other ranks, leaving each GPU with only its $1/N$ persistent shard of reduced gradients. Each worker then updates its $1/N$ parameter partition using the corresponding optimizer state, and an AllGather restores the updated parameter replicas. In the ideal ring accounting, the gradient ReduceScatter and the parameter AllGather each transfer $P(N-1)/N$ elements per worker, for $2P(N-1)/N$ across both phases, matching DDP's byte volume ($1\times$) when gradient and parameter communication widths match. Replicated parameters avoid forward parameter gathers, although model operations may still communicate. Under the FP16-gradient convention, persistent gradient storage falls from $2N$ bytes per parameter across N workers to 2 bytes per parameter in total, or $2/N$ per GPU, apart from transient bucket buffers.

The reductions in table 6.3 count memory per GPU rather than across the cluster. They are large-worker limits, in which the replicated FP16 weights and gradients remain on each GPU under ZeRO-1 and only the FP16 weights remain under ZeRO-2.

At the deepest stage, ZeRO-3 and FSDP partition the parameters as well, leaving each GPU with a $1/N$ persistent slice of all three training states. Because workers no longer hold full parameter

replicas between steps, the post-update parameter AllGather disappears and is replaced by two on-demand parameter AllGathers during computation. Each sharded unit gathers its parameters before the forward pass and immediately releases the non-local slices, then gathers them a second time before the backward pass and releases them again once local gradients are computed, and finally runs a ReduceScatter to deliver the completed gradients to each shard's owner for the local optimizer update. Across four ranks (figure 6.7) and under matching communication widths, ZeRO-1 ($M_{\text{optimizer}}/4$) and ZeRO-2 ($M_{\text{optimizer}}/4, M_{\text{gradients}}/4$) each execute two collective phases per step and retain the baseline $1\times$ communication volume. ZeRO-3 ($M_{\text{weights}}/4, M_{\text{gradients}}/4, M_{\text{optimizer}}/4$) executes three collective phases of one parameter-equivalent payload each (forward AllGather, backward AllGather, and gradient ReduceScatter), which yields $3/2 = 1.5\times$ the baseline two-phase volume in this accounting.

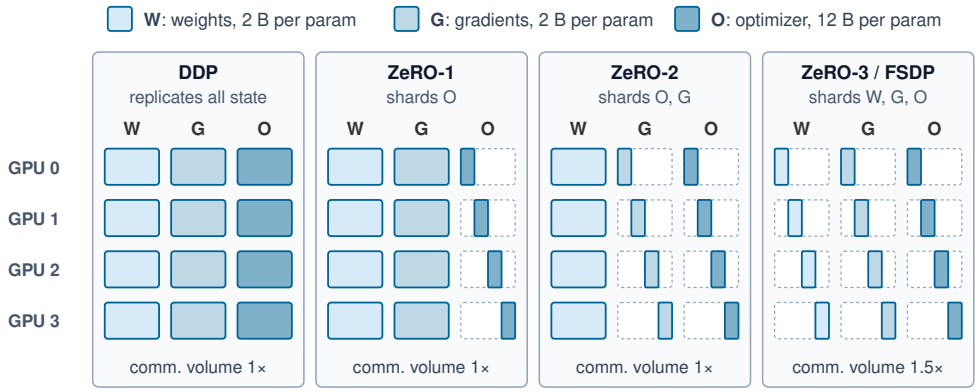

Figure 6.7: ZeRO State Partitioning across GPU Ranks: Which copy of each training state every GPU holds, for 4 workers. A filled full-width tile is a replicated copy; a quarter-width segment inside a dashed outline is that GPU's shard, stepping one quarter per rank. ZeRO-1 shards optimizer state ($M_{\text{optimizer}}/4$), ZeRO-2 adds gradients ($M_{\text{gradients}}/4$), and ZeRO-3 shards parameters too ($M_{\text{weights}}/4$). The footers give per-worker communication volume relative to DDP: only ZeRO-3 pays more, for the parameter AllGathers.

Full sharding adds parameter-readiness dependencies that replicated DDP avoids. With one sharding unit per layer, the forward pass needs an AllGather to reconstruct parameters (M_{layer} bytes). Resharding after forward requires another gather before backward (M_{layer} bytes), followed by a ReduceScatter for gradients (M_{layer} bytes under equal communication precision). For a model with N_L layers, this schedule performs about $3N_L$ collective operations per step. DDP instead AllReduces gradient buckets as they become ready. At large worker counts, the ideal per-worker transfer volumes approach $3M_{\text{param}}$ bytes versus $2M_{\text{param}}$ for DDP. Here M_{param} is the model's parameter bytes at the communication precision, not full training state, and it equals M_{grad} when parameters and gradients travel at the same width. These volumes do not determine elapsed time alone. While layer ℓ computes, layer $\ell + 1$ can prefetch parameters, subject to memory headroom and fabric contention.

The allocation lifetimes in figure 6.8 separate persistent ownership from the temporary buffers needed to execute one sharding unit.

A model fits only when all simultaneously live allocations fit. The persistent shards therefore establish a floor, while gathered units, saved activations, and prefetch determine how far the execution peak rises above it.

The choice between FSDP and DDP begins with peak memory. If replicated state, activations, and working buffers fit, DDP avoids repeated parameter gathers. When optimizer and gradient storage are the obstacle, ZeRO-2 offers savings while leaving parameters replicated. If parameters themselves do not fit, ZeRO-3/FSDP, offload, or partitioned computation must address that remaining constraint. For a 70B+ model on 80 GB devices, the need for tensor parallelism depends on the gathered layer size, activation budget, and throughput target; parameter count alone does not decide it.

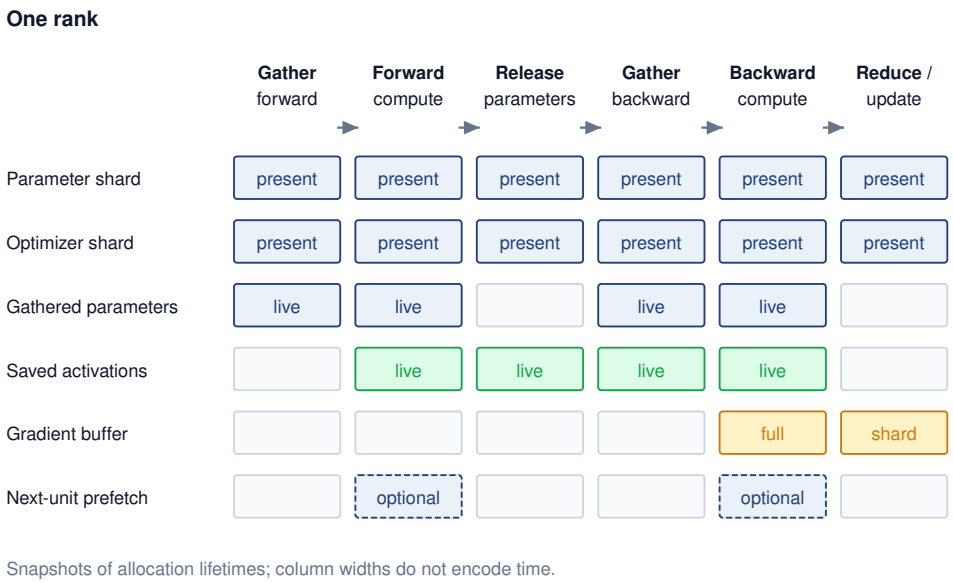

Figure 6.8: FSDP Parameter Lifetimes: One rank executes a sharding unit under a policy that reshards parameters after forward computation. Gathered parameters disappear between forward and backward, while saved activations remain until backward consumes them. Dashed cells show optional prefetch of another unit; these allocations can coexist with the current unit. The cells describe logical lifetimes, not byte totals or operation durations. Implementations may alias shard storage or reuse buffers.

Memory-efficient data parallelism requires careful tuning of sharding strategy (by layer, by transformer block, or flat) and mixed precision settings. Finer sharding reduces per-GPU memory but increases communication frequency, because more AllGather and ReduceScatter operations must execute per training step. Replicated or sharded, a data-parallel design earns its added workers only when they turn into faster progress toward target quality, and scaling efficiency measures that conversion.

6.4 Scaling Efficiency and Convergence

Data parallelism and ZeRO sharding show how communication and synchronization enter every training step, yet raw step throughput alone cannot judge a distributed run. When doubling the number of GPUs yields only 1.5× speedup for the same work, scaling efficiency is 75 percent, and if the larger global batch also requires more optimizer updates to reach the same loss, wall-clock time to quality degrades further. Evaluating a distributed configuration therefore requires separating two questions: how efficiently added devices execute each step, and how batch growth or stale updates affect the number of steps to convergence. The hardware scaling metrics below apply to every parallelism strategy; the convergence analysis that follows concerns data-parallel batch growth and synchronization slack, establishing a batch-size constraint that tensor and pipeline configurations must also respect when combined with data parallelism.

6.4.1 The mathematics of scaling efficiency

That first question is answered by scaling efficiency, the fleet-level ratio introduced with the Fleet Law in section 1.5.4 and reported as parallel efficiency in the GPT-2 scaling notebooks. If a model trains in time T_1 on one device, ideal (linear) scaling would train it in time T_1/N on N devices; communication overhead, pipeline bubbles, and load imbalance widen the gap between that ideal and the realized time T_N. Equation 6.1 formalizes the ratio:

$$\eta_{\text{scaling}} = \frac{T_1}{N \times T_N} \tag{6.1}$$

where T_N is the training time on N devices. Measured per step of the same work, T_1 is the quantity the Fleet Law writes as T_{compute}, the single-device compute time that its compute term divides by N. An efficiency of 1.0 means perfect linear scaling; an efficiency of 0.5 means the run achieves only half the ideal speedup.

Definition 6.4: Scaling efficiency

Scaling efficiency (η_{scaling}) is the ratio of actual training throughput to ideal linear throughput when increasing the number of ML compute devices (N).

1. **Significance**: The ratio in equation 6.1 measures added devices against the specified baseline. At 0.50 efficiency, a 10,000-GPU run delivers the throughput that 5,000 perfectly scaled devices would, so half of the accelerator-hours buy no progress relative to ideal scaling. The equivalence is to an idealized baseline. A real 5,000-GPU run would have its own, typically higher, efficiency, and the loss is spread across every device as waiting rather than concentrated in idle ones.
2. **Distinction**: Unlike single-node efficiency, which captures local bottlenecks like BW, scaling efficiency measures how performance changes with device count. Communication time $(T_{\text{comm}}(N))$ and synchronization contribute, but changing local kernel shapes or input delivery can also affect the ratio.
3. **Common pitfall**: A frequent misconception is that scaling efficiency is constant. In reality, it depends on both device count and problem size. For a fixed model, the communication-to-compute ratio typically worsens as N increases (Amdahl's Law). At a fixed N, a smaller model or local batch leaves less computation per device to hide startup latency and coordination, so its efficiency falls sooner.

The Fleet Law leads to the scaling efficiency metric for fixed global work. Substituting its step time (principle 10) for T_N in equation 6.1, with T_{compute} standing for T_1 per step as noted there, gives equation 6.2:

$$\eta_{\text{scaling}} = \frac{T_{\text{compute}}}{N \times T_{\text{step}}(N)} = \frac{1}{1 + \frac{N(T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}})}{T_{\text{compute}}}} \tag{6.2}$$

This form states the scaling efficiency bound (principle 12), under which perfect linear scaling $(\eta_{\text{scaling}} = 1.0)$ is a *theoretical limit*, not a *practical target*. Real systems fall short of it by an amount that depends on the configuration and grows with N. The 175B data-parallel example below reaches only 15 percent under its stated assumptions, and the hybrid layouts later in the chapter, priced on the reference training step rather than on this counterfactual, still stay below 1.0. The gap between $\eta_{\text{scaling}} = 1.0$ and the achieved efficiency is the price of distributed execution, paid in communication and coordination.

The equation also exposes the **scaling wall**. As N increases, the compute term (T_{compute}/N) shrinks, but the communication and synchronization terms can remain constant or grow. A ring AllReduce's bandwidth term, for example, stays near its bound however many workers join (principle 11). Eventually, the denominator is dominated by overhead, driving efficiency toward zero. The equation holds the global work fixed, which is strong scaling, and that regime reaches the wall sooner. Under weak scaling each device keeps its local batch, so its compute time per step stays level against the same ring term and the wall arrives later, while the growing global batch runs toward the statistical limit of section 6.4.6.

For data-parallel training of the 175B reference model, the per-step communication cost is dominated by the AllReduce of 350 GB of gradients, and how much of that exchange hides behind computation decides the efficiency. The scaling efficiency depends on computation per communicated byte. At fixed tokens per update, increasing parameter count raises both transformer computation and gradient bytes roughly proportionally, so model size alone does not guarantee better scaling. Larger batches raise computation per optimizer step without changing gradient size.

Large batches, however, can harm convergence and so require learning-rate tuning and warmup schedules.

Network bandwidth enters directly because doubling InfiniBand bandwidth halves the bandwidth term, but efficiency changes according to the full step-time denominator rather than in direct proportion. The network fabric is therefore a first-order determinant of cluster productivity, not a secondary concern. Every point of scaling efficiency it recovers is returned across the much larger accelerator investment, which is why fabric spending is judged by the efficiency it buys rather than by its own price.

These factors interact, so efficiency cannot be inferred from model size or GPU count alone. The following worked model fixes the parameter count, tokens per update, accelerator count, effective bandwidth, local compute efficiency before communication, and overlap before calculating one efficiency point. Changing any of those inputs requires recomputing the full step-time model.

Napkin Math 6.5: Scaling efficiency for a 175B model

Setup: A communication-only model for training a 175B model on a DGX H100 cluster with 400 Gb/s InfiniBand per GPU. Pure replicated data parallelism is a counterfactual baseline here: the calculation isolates gradient exchange but does not establish that the full training state fits on each device.

- **Compute per step** (assuming batch size 2M tokens, 6 FLOPs per parameter per token): $O_{\text{step}} = 6 \times 175 \times 10^9 \times 2 \times 10^6 \approx 2.1 \times 10^{18}$ FLOPs
- **Per-GPU compute time** on 1,024 GPUs, each at 1,979 TFLOP/s peak in the 8-bit floating-point (FP8) format developed in section 6.4.8 (50 percent local compute efficiency before communication). The general matrix multiplications (GEMMs) run in FP8 while the gradients below still cross the wire in FP16, because the collective's data type is chosen independently of the arithmetic's: $T_{\text{compute}}/N \approx 2.1$ seconds
- **AllReduce time** (for 350 GB of gradients using ring-AllReduce with overlap): $T_{\text{comm,exposed}}(N) \approx 11.9$ seconds (raw ring transfer takes $2 \times (N-1)/N \times 350/50$, about 14 seconds; the schedule could hide 75 percent of that transfer, but the available compute window caps hidden communication at 15 percent of raw transfer time)
- **Scaling efficiency**: $\eta_{\text{scaling}} = T_{\text{compute,local}}/(T_{\text{compute,local}} + T_{\text{comm,exposed}})$, which is equation 6.2 with the per-device compute window $T_{\text{compute,local}} = T_{\text{compute}}/N$ and the exposed term standing for $T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}}$. Here the numerator is 2.1 s and the complete denominator is 14 s, giving 15 percent. The exposed communication term already subtracts the hidden portion from the raw transfer time and includes collective synchronization; subtracting overlap again would double-count the same saving.

Systems insight: The low efficiency (15 percent) shows why naive data parallelism at this scale is insufficient. It is one operating point, fixed by 2M per update, 1,024 GPUs, and 15 percent overlap, not a universal threshold. Another cluster size requires recomputing per-device compute time, collective cost, overlap, and topology. A hierarchical reduction, which first reduces inside each node over NVLink so that each InfiniBand port carries only a fraction of the gradient, shrinks the exposed term, and so does a longer compute window per step. Section 6.6 prices both on the reference training step, where memory rather than communication rules out pure data parallelism.

↩ The compute window that bounds how much of a collective can hide is established in section 5.10.4.

6.4.1.1 Parallelism-infrastructure interaction

The scaling efficiency of a run is set by the *interaction* between the model's communication requirements and the infrastructure's bandwidth hierarchy, not by the *model architecture alone*. Making a large model both fit and run efficiently means mapping each parallelism dimension onto the bandwidth tier that can carry its traffic. Section 6.6 works through four concrete configurations, priced on the reference training step, once tensor and pipeline parallelism have been developed.

The collective algorithm adds its own dependence on message size. Tree-based AllReduce reaches every rank in $\mathcal{O}(\log N)$ steps and suits small, latency-bound messages, whereas large gradient buffers favor the bandwidth-optimal ring; message size, topology, and contention decide between them.

Even pre-frontier transformers carry gradient buffers large enough to demand substantial link bandwidth from that ring. BERT-Base (110M parameters) requires approximately 440 MB of gradient synchronization per iteration in FP32, while BERT-Large (340M parameters) requires approximately 1.4 GB. Because a ring AllReduce across 64 GPUs moves nearly twice the gradient bytes through each GPU's link, finishing BERT-Large synchronization within 50 ms demands roughly 54 GB/s of sustained per-GPU bandwidth, and BERT-Base about a third of that. At 175B parameters, meeting the step budget requires combining 3D parallelism, gradient accumulation, and overlap rather than relying on a single flat collective, and whatever communication remains unhidden acts as the unscalable residual that Amdahl's Law prices next.

↩ The message size at which ring AllReduce overtakes tree AllReduce, derived from the latency and bandwidth terms of each algorithm, is established in section 5.6.7.

6.4.2 Scaling physics and Amdahl's Law with communication

Even a small unscalable cost in a single-device step eventually dominates distributed execution, because adding N devices shrinks the parallelizable compute share (T_1/N) while leaving the fixed residual untouched. An Amdahl-style model separates that scalable work from the fixed residual fraction r of the single-device baseline (section C.3.2). In units of the single-device step, a distributed step then takes the scalable share divided by the device count plus the whole residual, and speedup is the reciprocal of that time.

That baseline fraction is not the synchronization share a profiler reports for a distributed step (the distributed Amdahl pitfall of section 1.5.4), because the compute share has been divided by N while the residual has not. A residual of 1 percent of the single-device step occupies about 84 percent of each step on 512 GPUs. Reading the observed share as r would therefore overstate the residual many times over and predict a far lower speedup ceiling than the model implies.

Plotted on log-log axes up to 512 GPUs, the same fixed-residual model shows how much speedup each residual fraction gives up (figure 6.9), complementing the scaling efficiency bound (principle 12). In this illustrative model, a small baseline fraction ($r = 0.1\%$) stays closer to ideal scaling than a larger fraction ($r = 1\%$). A residual equal to a tenth of baseline time ($r = 0.10$) limits asymptotic speedup to $1/r = 10\times$. Actual collective and coordination costs can change with worker count, so predicting a training run requires the full step-time accounting of equation 6.2.

Published benchmark runs show the same shortfall in practice. MLPerf Training times complete runs to a fixed quality target, and several submitters report the same system at several sizes within one round (Mattson et al. 2020; MLCommons 2025b). Each of 5 such series for the GPT-3 and Llama 3.1 405B benchmarks falls below the ideal line drawn through its smallest configuration (figure 6.10). On NVIDIA's Eos cluster, GPT-3 trained in about 58 minutes on 512 H100 GPUs and 3.9 minutes on 10,752. A 21× increase in accelerators therefore cut the time 14.9×, a scaling efficiency (equation 6.1, measured relative to the smallest configuration rather than one device) of 71 percent. Across the series, the largest configuration keeps 71–84 percent of ideal. Two effects compound in these numbers. Communication and coordination grow with worker count, and the larger configurations also train with larger global batches, which can require more samples to reach the target (section 6.4.6). Separating the two requires the step-level accounting of equation 6.2.

↪ The single-machine iron law and its data-movement term are developed in section A.2.

Predicting a run means pricing each term of equation 6.2 for the configuration at hand. For a large-message ring AllReduce, the communication term's bandwidth contribution is $\frac{2(N-1)}{N}\frac{D_{\text{vol}}}{\beta}$. This is the fleet counterpart of the single-machine iron law's data-movement ratio D_{vol}/BW. Here D_{vol} is the payload being reduced (the gradient bytes M_{grad} under data parallelism, or a shard of model state under sharding), and β is sustained per-link bandwidth. The expression omits per-message startup latency, and a wait it already counts must not be charged again under synchronization.

This communication overhead costs energy as well as wall-clock time, and the energy cost of communication grows with the physical distance between devices. At the scale of 10,000 GPUs, moving data is itself a power load. A full H100 NVLink envelope of roughly 9 PB/s at 7.5 pJ/bit represents about 540 kW before cooling and switch overhead. The smaller NDR InfiniBand envelope (500 TB/s at 35 pJ/bit) still represents about 140 kW. The per-boundary accounting behind these

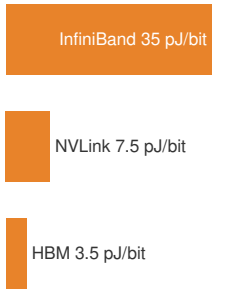

Communication energy climbs from HBM out to the network.

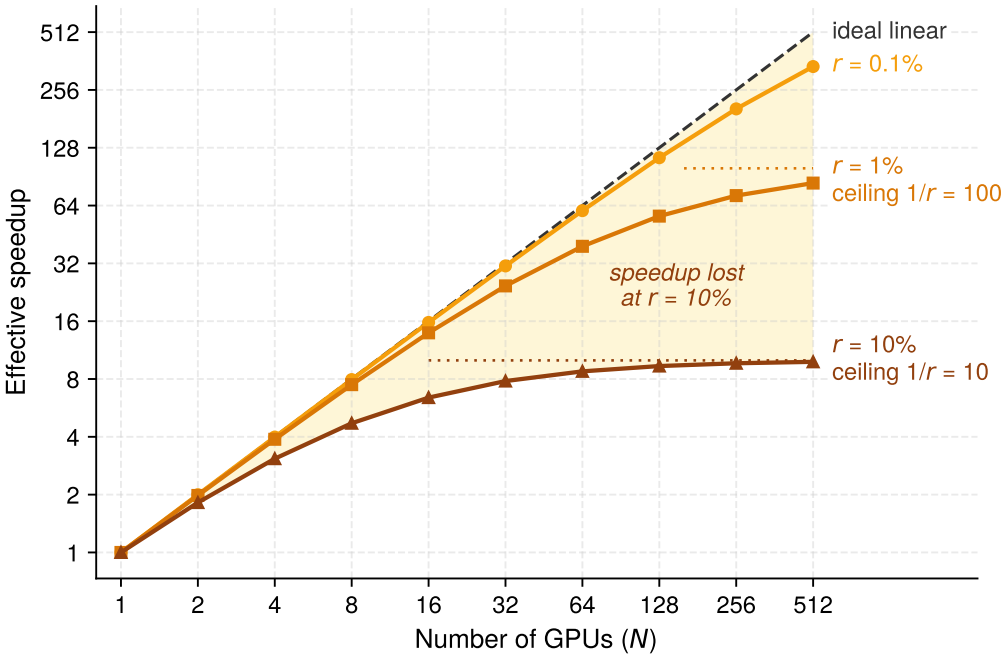

Figure 6.9: Speedup under a Fixed Residual: Effective speedup $N/(1+(N-1)r)$ against GPU count on log-log axes, for illustrative fixed residual fractions r of baseline step time. Ideal scaling (dashed) is linear. Each curve bends away from it and flattens near $1/r$: at 512 GPUs, $r = 0.1$ percent still delivers about two-thirds of ideal speedup, 1 percent about one-sixth, and 10 percent under $10\times$ in total. Darker curves carry larger r.

per-bit costs is developed in section 5.2. Overlap hides communication time but not its energy, so avoiding unnecessary movement is the direct way to reduce the power term.

As the device count grows for a fixed workload, the balance can move from a linear scaling regime toward a communication-bound regime. The transition has no universal GPU count, because it depends on compute per device, message sizes, topology, and overlap. Coordination overhead becomes consequential when its exposed time is no longer small relative to useful computation. Measurements at the intended workload and placement must establish the efficiency curve.

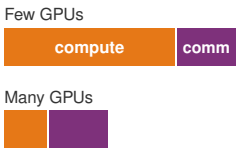

More GPUs shrink each step's compute, not its communication.

The efficiency metrics directly influence the choice of parallelism strategy. Data parallelism works well in the linear scaling regime but becomes communication bound at scale. The model-parallel strategies developed after this section replace part of that gradient traffic with other costs, which the same efficiency equation prices once each mechanism is in place.

Hardware efficiency metrics govern throughput, yet convergence theory determines whether distributed training reaches the same solution quality as single-device training. Parallelism affects optimization convergence in four ways: convergence rate changes with batch size, stale updates can weaken that rate, learning-rate schedules must be revalidated when the global batch changes, and adding workers yields diminishing returns beyond the critical batch size.

6.4.3 Convergence rate for synchronous data parallel SGD

Averaging gradients across N workers processes N times as many examples per step, but only reduces the number of steps to convergence when gradient noise, rather than loss curvature, is what holds back a larger step size. Formally, let $\mathcal{L}$ be lower bounded by $\mathcal{L}_{\text{inf}}$ and have L_s-Lipschitz gradients, so the smoothness constant limits how sharply the gradient can change as parameters move. Each worker computes gradients from B_{local} independent samples. At update t, the aggregated estimator $\mathbf{g}_t$, the average of the N workers' local gradients, is unbiased, and its conditional variance satisfies $\mathbb{E}[\|\mathbf{g}_t - \nabla\mathcal{L}(\theta_t)\|^2 \mid \theta_t] \le \sigma^2/(NB_{\text{local}})$. When averaging reduces that noise, synchronous SGD makes steady progress toward parameters where the gradient is small. These assumptions

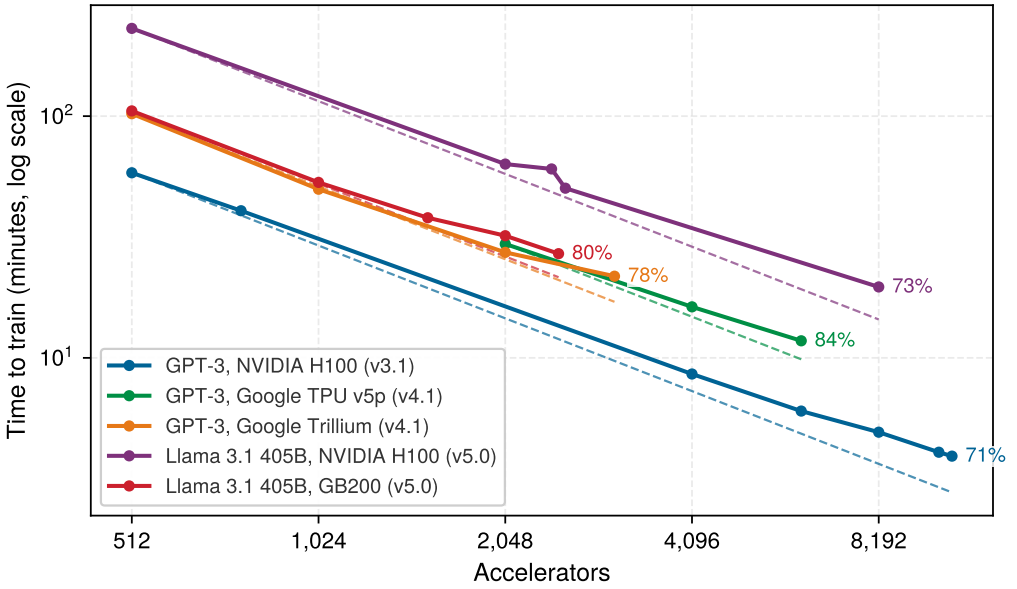

Figure 6.10: Measured Scaling in MLPerf Training: Time to train against accelerator count on logarithmic axes for five same-round series of the GPT-3 175B and Llama 3.1 405B benchmarks. Each dashed line is ideal scaling from that series' smallest configuration, and the label gives the scaling efficiency of its largest configuration relative to its smallest. Global batch size grows with scale in several series, so the gap from ideal includes batch-dependent convergence as well as communication and coordination. Times are recomputed from the published run logs and can differ slightly from official scores. Data from MLCommons (2025b).

bound progress toward a small gradient norm, not convergence to a global minimizer or a particular validation accuracy.

Theorem 6.1: Convergence rate for distributed SGD

For the update $\theta_{t+1} = \theta_t - \eta \mathbf{g}_t$ with constant step size $0 < \eta \leq 1/L_s$, after T synchronized updates,

$$\frac{1}{T}\sum_{t=0}^{T-1} \mathbb{E}\left[\|\nabla \mathcal{L}(\theta_t)\|^2\right] \leq \underbrace{\frac{2(\mathcal{L}(\theta_0) - \mathcal{L}_{\text{inf}})}{\eta T}}_{\text{optimization term}} + \underbrace{\frac{\eta L_s \sigma^2}{N B_{\text{local}}}}_{\text{stochastic-variance term}} .$$

Choosing $\eta = \min\left\{1/L_s, \sqrt{2(\mathcal{L}(\theta_0) - \mathcal{L}_{\text{inf}}) N B_{\text{local}}/(L_s \sigma^2 T)}\right\}$ balances the two terms. Before the smoothness cap becomes active, the stochastic contribution scales as $\mathcal{O}(1/\sqrt{N B_{\text{local}} T})$.

The theorem isolates the benefit that follows from the assumptions. With equal local batch size B_{local}, synchronous data-parallel SGD computes the same minibatch estimator as one worker using batch size $N \times B_{\text{local}}$ for that update, and the estimator variance decreases as $1/(N \times B_{\text{local}})$. That product appears only in the stochastic-variance term of the bound, so at a fixed step size a larger batch shrinks that term and leaves the optimization term alone. The reduction in updates comes through the step size. A smaller variance term lets the balanced step size grow, and the updates the bound requires fall in inverse proportion to the batch. Once the step size sits at the smoothness cap, the optimization term sets the update count, and added samples no longer reduce it in proportion. The critical batch size marks that diminishing return. How far added samples reduce those updates is the statistical efficiency of distributed training, separate from hardware efficiency. Because the theorem assumes perfect synchronization (BSP), convergence guarantees degrade as staleness increases whenever workers proceed at different rates or apply stale gradients.

6.4.4 Staleness impact across synchronization models

When a slow worker applies a gradient computed on an older parameter version, the update pushes the model along a direction that the current parameters may have already outgrown. The staleness parameter τ_{stale} counts how many updates occur between computing and applying a gradient, tying the barrier savings of BSP, SSP, and ASP (section 6.3.1.5) to solution quality (Ho et al. 2013; Dutta et al. 2018). Because the damage from stale gradients depends on the objective, step-size schedule, worker count, and delay distribution, no single additive penalty applies across all runs.

Definition 6.5: Gradient staleness

Gradient staleness (τ_{stale}) is the number of parameter updates that occur between the time a gradient is computed and the time it is applied to the global model state.

1. **Significance**: It represents the synchronization error in distributed optimization. Increasing τ_{stale} can improve throughput by reducing barrier waits ($T_{\text{sync}}(N)$), but it typically degrades the rate of convergence, requiring more optimizer steps, and therefore more FLOPs, to reach the same accuracy.
2. **Distinction**: Unlike network latency, which is a physical delay, staleness is an algorithmic offset that arises from the choice of synchronization protocol (e.g., ASP, SSP).
3. **Common pitfall**: A frequent misconception is that staleness is "always bad." In reality, it is a throughput-convergence trade-off: for some large-scale workloads, allowing bounded staleness is one practical way to keep thousands of GPUs in use.

The convergence behavior differs across these models in ways that directly affect training cost and solution quality. In BSP ($\tau_{\text{stale}} = 0$), all workers compute on the same parameter version and aggregate before updating. Under the assumptions of theorem 6.1, the synchronized estimator inherits the $N \times B_{\text{local}}$ minibatch guarantee.

SSP ($\tau_{\text{stale}} \leq s$) allows workers to proceed up to s iterations ahead of the slowest worker. Ho et al. (2013) bound the average regret of their SSP update sequence at $\mathcal{O}(\sqrt{(s+1)N/T})$ for N workers and T updates. Average regret is the per-update gap between the loss its iterates incur and the loss of the best fixed parameters chosen in hindsight. The bound assumes convex Lipschitz component functions, a bounded domain, and a decreasing step size. It makes the staleness dependence explicit, since the bound grows with s. It measures a different quantity from the gradient-norm bound of theorem 6.1, which covers nonconvex objectives, so the two rates cannot be compared directly.

Asynchronous SGD imposes no protocol-wide finite delay bound. Removing the barrier can improve update throughput, but a convergence rate requires additional assumptions about realized delays, the objective, and the step-size schedule. Dutta et al. (2018) analyze this as an error-runtime trade-off in which worker speed and gradient staleness jointly determine whether accepting a delayed update helps. Asynchronous execution alone therefore sets no fixed delay penalty and does not by itself forfeit the N-worker variance reduction; both depend on those assumptions. The protocol fixes each model's delay control, but the strength of its convergence guarantee depends on further assumptions (table 6.4).

Table 6.4: Convergence Properties by Synchronization Model: The synchronization protocol determines the delay model; the convergence rate follows only after the objective, gradient, delay, and step-size assumptions are fixed.

Model	Delay control	Scoped guarantee	Systems effect
BSP	$\tau_{\text{stale}} = 0$	Synchronous $N \times B_{\text{local}}$-minibatch theorem 6.1	Waits for the slowest worker
SSP	$\tau_{\text{stale}} \leq s$	$\mathcal{O}(\sqrt{(s+1)N/T})$ average regret under convex assumptions (Ho et al. 2013)	Bounds delay while reducing some waits
ASP	No protocol-wide finite bound	No universal rate without additional delay and objective assumptions	Removes the global barrier but applies stale updates

6.4.5 Learning rate scaling rules

Even under zero-staleness BSP ($\tau_{\text{stale}} = 0$ in table 6.4), scaling out across N workers changes the optimization trajectory whenever the global batch grows with the cluster. Multiplying the global batch by k cuts the number of optimizer updates per epoch by the same factor, so keeping the single-GPU learning rate leaves the model under-trained, while jumping straight to a k-fold larger step size can destabilize early epochs. Whenever data parallelism changes the effective batch size, the learning-rate schedule must be revalidated against the target quality; two scaling rules provide starting points for that adjustment.

The linear scaling rule (Goyal et al. 2017) proposes that when batch size is multiplied by k, learning rate also increases by k:

$$\eta_{\text{new}} = k \cdot \eta_{\text{base}}$$

The rule rests on an approximation that compares one large-batch step at learning rate $k\eta$ with k consecutive small-batch steps at learning rate η. At batch size $k \times B$, independent gradient contributions reduce estimator variance by factor k, but lower variance alone does not establish stability. The approximation also requires gradients to change little over the corresponding k small-batch steps, because only then do those steps add up to nearly the same update as the one large step. Goyal et al. (2017) validated the recipe for ResNet-50/ImageNet at global batch size 8,192 with gradual warmup. A warmup period increases the learning rate from η_{base} to $k \cdot \eta_{\text{base}}$ over the first W iterations to reduce early instability; the schedule still requires validation:

$$\eta_t = \eta_{\text{base}} + \frac{t}{W}(k \cdot \eta_{\text{base}} - \eta_{\text{base}}) \quad \text{for } t < W$$

The rule and the warmup were introduced together, and the run that introduced them shows why the scaled learning rate needs the warmup. A learning-rate schedule tuned for an eight-GPU run need not reach the same accuracy at 256 GPUs once the global batch grows with the worker count.

War Story 6.1: The one-hour ImageNet run (2017)

Context: Facebook AI Research set out to train ResNet-50 on ImageNet in one hour using 256 GPUs, pushing data parallelism far beyond the batch sizes that earlier practice treated as safe (Goyal et al. 2017).

Mechanism: Scaling the global batch to 8,192 cut the number of optimizer updates per epoch by the same factor. Keeping the small-batch learning rate left too few effective updates to reach the target accuracy, while scaling the learning rate up by the same factor destabilized optimization in the first epochs, when the loss surface changes fastest.

Impact: Without the gradual warmup recipe, the large-batch runs reported higher training and validation error than the small-batch baseline. Higher hardware throughput alone did not preserve the quality target.

Fix: The team introduced the linear learning-rate scaling rule paired with a 5-epoch gradual warmup schedule, stabilizing optimization at batch size 8,192.

Systems lesson: Distributed training demands more than parallel hardware alone. Scaling changes the optimization regime, so the cluster, communication schedule, batch size, and learning-rate schedule must be tuned as one system.

The square root scaling rule is an alternative learning-rate heuristic to test when increasing batch size:

$$\eta_{\text{new}} = \sqrt{k} \cdot \eta_{\text{base}}$$

This rule increases the learning rate more slowly than linear scaling. Its suitability depends on the optimizer, schedule, and training regime; exceeding a critical-batch estimate does not make it mandatory or guarantee convergence. The learning rate must be validated against the same target quality.

Layer-wise adaptive rate scaling (LARS) and layer-wise adaptive moments based optimizer (LAMB) provide related large-batch approaches for ImageNet and BERT-style pretraining (You et al. 2017; You et al. 2020). LARS (You et al. 2017) scales each layer's update relative to its weight norm. The simplified LARS expression below illustrates the ratio of weight norm to gradient norm with weight decay. LAMB (You et al. 2020) applies a related layer-wise adaptation to an Adam-based update rather than using this raw-gradient denominator:

$$\eta_\ell = \eta_{\text{global}} \cdot \frac{\|\mathbf{w}_\ell\|}{\|\mathbf{g}_\ell\| + \lambda_{\text{reg}}\|\mathbf{w}_\ell\|}$$

Here λ_{reg} is the weight-decay coefficient in the illustrated LARS expression. Layer-wise scaling controls update magnitude relative to parameter magnitude, but does not guarantee convergence for arbitrary batches. Its systems role is to extend the range of validated batch sizes over which added data-parallel workers can reduce time to quality.

6.4.6 Critical batch size and diminishing returns

Increasing the global batch eventually yields diminishing reductions in the optimization steps needed to reach a fixed target quality. The critical batch size B^* marks the transition point beyond which increasing batch size yields diminishing returns in convergence per sample seen.

Definition 6.6: Critical batch size

Critical batch size (B^*) is a workload-dependent crossover batch size beyond which further batch growth yields diminishing reductions in optimization steps per added sample (McCandlish et al. 2018; Shallue et al. 2019).

1. **Significance**: It characterizes statistical scaling, separately from hardware throughput. In the gradient-noise model, batches well below B^* can reduce step count approximately in proportion to batch growth; above it, additional samples buy progressively smaller step-count reductions.
2. **Distinction**: Unlike the memory-limited batch size (determined by memory capacity and activation footprint), the critical batch size depends on the model, data, optimizer, training phase, and target quality.
3. **Common pitfall**: A frequent misconception is that training can be sped up indefinitely by adding GPUs. In reality, B^* is a guide to diminishing statistical returns, not a hard physical ceiling or a universal optimum in wall-clock time.

The gradient-noise-scale proxy for the critical batch size can be estimated as:

$$B^* \approx \frac{\text{tr}(\Sigma)}{\|\nabla \mathcal{L}(\theta)\|^2}$$

where $\text{tr}(\Sigma)$ is the trace of the gradient covariance matrix (total gradient variance) and $\|\nabla \mathcal{L}(\theta)\|^2$ is the squared gradient norm (signal strength). Intuitively, B^* is the batch size at which averaging reduces gradient variance to the level of the true gradient magnitude.

Published large-batch regimes illustrate the scale rather than providing fixed thresholds. On ImageNet, Goyal et al. (2017) stabilized ResNet-50 at global batch 8,192 with warmup, and LARS extends ImageNet large-batch training further in some settings (You et al. 2017). For BERT-Large pretraining, LAMB reports training at 32K–64K batch sizes with layer-wise adaptive updates (You et al. 2020). Across other domains, McCandlish et al. (2018) show useful batch-size limits ranging from tens of thousands in ImageNet-like settings to millions in reinforcement-learning workloads, so new frontier-language-model thresholds should be measured rather than copied.

The gradient-noise model separates statistical efficiency from hardware efficiency. Well below B^*, where gradient noise rather than loss curvature holds back the step size, increasing batch size can reduce optimization steps almost proportionally while preserving similar total sample use. Near B^*, the trade-off between fewer updates and more samples becomes material. Above it, curvature caps

the step size, so step count approaches a lower limit while sample use continues to increase. Neither maximum samples per second nor minimum training time necessarily occurs at B^*: those objectives also depend on device utilization, communication, and the validated learning-rate schedule.

On log-log axes, the gradient-noise trade-off appears as a divergence between optimization step count and total sample consumption (figure 6.11). In the normalized model plotted there, steps relative to their minimum equal one plus the critical batch divided by the batch, and samples relative to their minimum equal one plus the batch divided by the critical batch. Below the critical threshold ($B/B^* < 1$), increasing batch size yields near-linear reductions in optimization steps with negligible sample waste. Past the crossover point at $B/B^* = 1$ (where both normalized costs equal 2), the step count curve asymptotically flattens toward 1 while total examples processed climb steeply along the $1 + B/B^*$ trajectory.

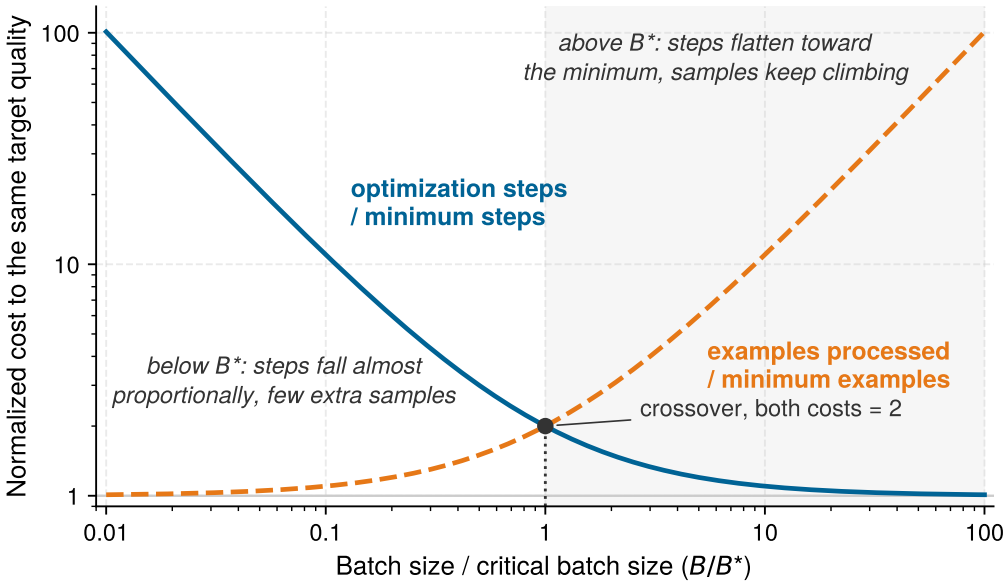

Figure 6.11: Statistical Cost of Larger Batches: A normalized illustration of the gradient-noise trade-off model, with batch size divided by the fitted critical batch. Optimization steps approach a lower limit as batches grow, while total examples processed increase. The crossover is not a hardware-throughput optimum or a measured threshold for a particular model (McCandlish et al. 2018).

Dividing the critical batch by the per-worker batch B_{local} estimates where worker growth leaves the near-proportional statistical scaling regime. With one local batch per update and its size held fixed, the crossover occurs around $N \approx B^*/B_{\text{local}}$ workers. Past that point, additional workers mainly raise sample throughput that the optimizer cannot convert into proportionally fewer steps. Gradient accumulation increases each worker's contribution per update and lowers the corresponding worker count. Shrinking B_{local} moves the crossover outward but leaves less computation per worker to hide communication, coupling statistical and hardware efficiency.

That crossover informs a cost-time trade-off rather than setting a mandatory worker count. Larger batches may still reduce time to target even as they consume more examples. That benefit must be established with measured step time and a validated learning-rate schedule at the same target quality. Changing synchronization or compression also changes the optimization process, so its communication savings must be weighed against any additional steps required.

Checkpoint 6.2: Scaling decisions

Given a 7B-parameter model distributed across a cluster of 64 A100 GPUs (80 GB HBM2e each), estimate the critical-batch crossover. Use a pilot trace with $\text{tr}(\Sigma) = 1.6 \times 10^6$, $\|\nabla \mathcal{L}(\theta)\|^2 = 100$, per-GPU batch $B_{\text{local}} = 4$, and a proposed gradient-accumulation depth of 32 steps, giving $B_{\text{global}} = 64 \times 4 \times 32 = 8{,}192$ samples.

- ☐ **Critical batch size**: calculate the critical batch size (B^*), where the gradient noise scale equals the batch size.
- ☐ **Proposed batch**: use the critical batch size proxy $B^* \approx \frac{\text{tr}(\Sigma)}{\|\nabla \mathcal{L}(\theta)\|^2}$ to evaluate the proposed batch size.
- ☐ **Efficiency distinction**: explain why the gradient-noise estimate can assess statistical scaling but cannot establish hardware scaling efficiency without compute and communication timings.

6.4.7 Communication cost versus convergence speed

Communication savings matter only through the resulting time and cost to reach target model quality. Scaling a transformer language model from 8 to 64 workers, with a per-worker batch of $B_{\text{local}} = 32$, shows how shorter steps and extra iterations to the same target combine into that cost.

Napkin Math 6.6: Scaling from 8 to 64 workers

Setup: This illustrative comparison assumes the iteration counts needed to reach one fixed target quality. They are scenario inputs, not measured convergence results or predictions from GPU count. The baseline is 100K iterations on single GPU with $B_{\text{local}} = 32$.

8 workers (BSP)

- Effective batch size: $B = 8 \times 32 = 256$
- Learning rate: $\eta = 8 \times \eta_{\text{base}}$ (illustrative candidate with warmup; validate for the workload)
- Expected iterations: 100K / 8 = 12.5K iterations
- Assumed iterations to the same target: 12.8K iterations (98 percent efficiency)
- Communication overhead: 15 percent (NVLink intra-node)
- Wall-clock speedup: 100K × 1 / (12.8K × 1.15×) = 6.8×

64 workers (BSP)

- Effective batch size: $B = 64 \times 32 = 2{,}048$
- Learning rate ($N = 64$ workers): $\eta = N \times \eta_{\text{base}}$ is the scenario's assumed candidate schedule, not a prescription determined by B^* alone
- Assuming $B^* \approx 4{,}000$ (below critical): the useful-worker bound is $B^*/B_{\text{local}} \approx 4{,}000 \ / \ 32 \approx 125$ workers, so 64 workers stay inside the near-proportional regime; this estimate does not validate a learning-rate schedule
- Expected iterations: 100K / 64 = 1.56K iterations
- Assumed iterations to the same target: 1.72K iterations (91 percent efficiency)
- Communication overhead: 45 percent (InfiniBand inter-node, 8 nodes)
- Wall-clock speedup: 100K $\times$ 1 / (1.72K $\times$ 1.45×) = 40.1×

SSP workers: 64 workers ($s = 4$)

- Same effective batch size: $B = 2{,}048$
- Illustrative learning-rate reduction, not a staleness theorem: $\eta' = \eta_{\text{BSP}}/\sqrt{1+s}$ with $s = 4$, giving $\eta' \approx 0.45 \times \eta_{\text{BSP}}$
- Expected iterations: Higher due to staleness penalty
- Assumed iterations to the same target: 2.1K iterations (74 percent efficiency)
- Communication overhead: 25 percent (reduced synchronization)
- Wall-clock speedup: 100K $\times$ 1 / (2.1K $\times$ 1.25×) = 38.1×

Analysis (table 6.5)

Table 6.5: Illustrative Scaling Comparison: Assumed iteration counts to one target quality, communication increments relative to local compute time, and the resulting speedup and sample efficiency. These are scenario calculations, not empirical optimizer benchmarks.

Configuration	Iterations	Comm. overhead	Wall-clock speedup	Sample efficiency
1 GPU (baseline)	100,000	0%	1×	100%
8 GPU BSP	12,800	15%	6.8×	98%
64 GPU BSP	1,720	45%	40.1×	91%
64 GPU SSP	2,100	25%	38.1×	74%

The 64-GPU BSP configuration achieves 40× speedup despite only 91 percent sample efficiency because the communication overhead (45 percent) is offset by the larger worker count. SSP at 64 workers comes close in wall-clock time with lower per-iteration communication (25 percent), but the extra iterations its stale updates require hold it to 38.1× against BSP's 40.1×. In this scenario it therefore finishes slightly later and costs slightly more than 64-worker BSP.

Cost Analysis (assuming $2/GPU-hour):

- 8: 12.8K iters × 0.46 s/iter × 8/3600 × $2/GPU-hour = $26
- 64 BSP: 1.72K iters × 0.58 s/iter × 64/3600 × $2/GPU-hour = $35
- 64 SSP: 2.1K iters × 0.50 s/iter × 64/3600 × $2/GPU-hour = $37

Despite higher parallelism, 64-GPU training costs more per run due to communication overhead and reduced sample efficiency. The 8-GPU configuration is more cost-efficient, but its 6.8× speedup against the 64-GPU run's 40.1× leaves it several times slower in wall-clock time. The choice depends on whether minimizing cost or minimizing time-to-result is the priority.

On the cost-time plane of figure 6.12, 64-worker BSP dominates 64-worker SSP on both axes. SSP saves communication per iteration, but its stale updates need enough extra iterations to cost more and finish later, while 8-worker BSP remains the cheaper and slower option.

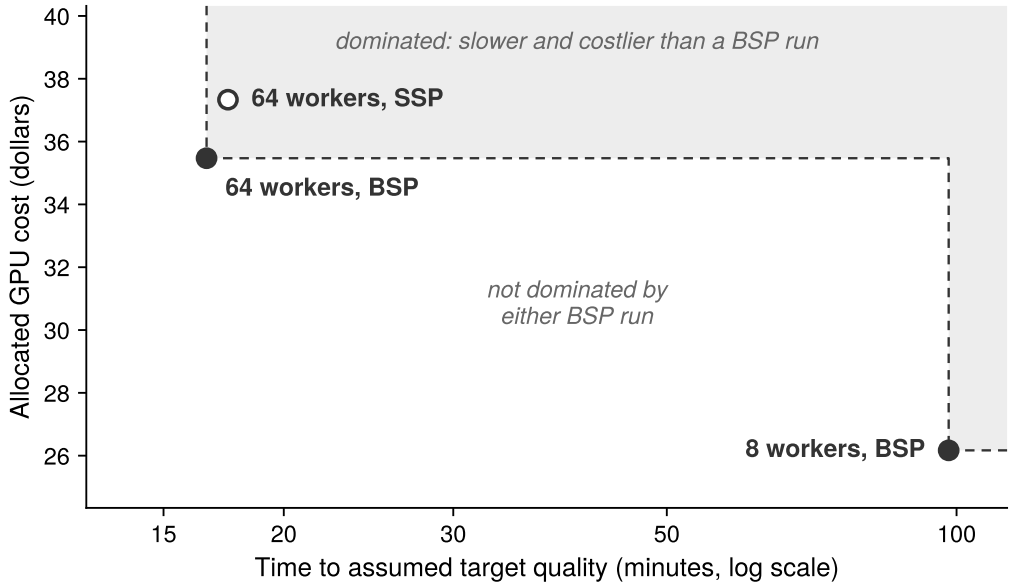

Figure 6.12: Cost and Time in the Illustrative Scaling Scenario: Points reuse the assumed convergence counts and step times in the preceding notebook. The eight-worker configuration costs less but finishes later; the assumed sixty-four-worker SSP configuration is dominated by BSP in both cost and time. This is not a universal ranking of synchronization methods.

Moving a job along this cost-time frontier requires shrinking each message, synchronizing less often, or relaxing the global reduction topology. Gradient compression shrinks the payload directly by transmitting quantized, sparsified, or low-rank approximations. Examples include sign-based updates (Bernstein et al. 2018), Deep Gradient Compression (Lin et al. 2018), PowerSGD (Vogels et al. 2019), quantized SGD with convergence analysis (Alistarh et al. 2017), and empirical sparse gradient dropping for neural machine translation (Aji and Heafield 2017).

↩ Quantized, sparsified, and low-rank gradient compression, the error feedback that keeps them convergent, and the criteria for when they pay, are established in section 5.9.

Instead of shrinking each payload, Local SGD lowers synchronization frequency by letting workers perform H_{local} local updates before averaging, which cuts collective calls per local update by a factor of H_{local}. Results for smooth, strongly convex objectives can recover synchronous SGD's asymptotic rate under specified learning-rate and local-update-interval conditions (Stich 2019), so

the synchronization interval remains part of the convergence contract rather than an unrestricted communication saving.

Where Local SGD keeps a global AllReduce at wider intervals, Decentralized SGD replaces the global collective with sparse neighbor exchanges along a communication graph. Immediate bandwidth demand drops, at the cost of taking additional rounds to spread information across the group. Graph connectivity, link service time, and the optimization effect of delayed mixing determine whether that trade-off helps a geographically distributed workload.

Choosing among these three levers depends on which bottleneck dominates and how many extra iterations the altered update costs. Each candidate must be judged by measured time to the same target quality. A narrower numeric format offers one more way to shrink the payload without changing the collective schedule, provided the collective transmits the narrower values.

6.4.8 FP8 for distributed training

Running matrix multiplications in 8-bit floating point (FP8) accelerates local compute, which can make a distributed job *more* communication-bound if gradients still cross the network in 16-bit precision. Compute precision and communication precision are separate choices: casting GEMM operands to half as many bytes reduces operand storage and local memory traffic, but it halves a collective payload only when the communication path explicitly transmits that narrower representation.

FP8 matters to distributed training only after four dtypes are distinguished: the stored parameter dtype, the GEMM operand dtype, the accumulator or master-state dtype, and the collective wire dtype. Traditional mixed-precision training commonly retains FP32 optimizer or master state while using FP16/BF16 operands. FP8-capable accelerators add E4M3 (four exponent bits and three mantissa bits) and E5M2 (five exponent bits and two mantissa bits) operand formats optimized for different tensor distributions (Micikevicius et al. 2022).

↩ The throughput that narrower precisions buy on current accelerators is established in section 2.8.1.

The two FP8 formats trade precision against range. E4M3 keeps an extra mantissa bit for scaled *weights and activations*, which cluster within a predictable range, while E5M2 keeps an extra exponent bit for scaled *gradients*, which can span many orders of magnitude during backpropagation. Because tensor magnitudes drift during training, each tensor is multiplied before conversion by a scale, commonly derived from its recent maximum absolute value, a policy called *dynamic scaling*. Neither the format nor the scaling policy guarantees that every value survives conversion, so training quality must be validated for each workload.

An FP8 training contract pairs FP32 master or optimizer state and FP8 GEMM operands with higher-precision accumulation, and it can approach FP16 training quality on suitable workloads (Micikevicius et al. 2022). FP8 halves bytes only for tensors actually stored or transmitted in FP8, so a collective that reduces FP16 gradients still moves FP16 bytes even when every preceding GEMM ran in FP8. FP8 therefore earns its place in distributed training when lower-precision compute shortens the local critical path, or when an explicit communication cast shrinks payloads while the accumulation and scaling policy protects convergence. It changes representation, not the parallelization axis. Compression, local updates, and narrower formats all make each data-parallel step cheaper, yet none of them partitions an individual layer or its activations; that remaining capacity constraint motivates model parallelism.

↪ Per-parameter training memory under each of these formats is collected for reference in section A.6.

6.5 Model Parallelism

Sharded data parallelism distributes stored state but still asks each worker to execute every layer on its local examples. If one layer's gathered parameters and working memory exceed device capacity, or its execution is too slow for the target schedule, splitting that computation becomes a candidate. Model parallelism addresses this form of the memory capacity gap (principle 2) by assigning different computations to different devices. Its value must be compared with offload, recomputation, and smaller local batches under the same throughput and quality requirements.

Definition 6.7: Model parallelism

Model parallelism is a distributed training strategy that partitions a single neural network's parameters or operations across multiple devices so each device computes a distinct portion of the model.

1. **Significance**: If a model with parameter state M_{weights} is split across $N_{\text{TP}}N_{\text{PP}}$ model-parallel devices, the ideal per-device parameter footprint falls toward $M_{\text{weights}}/(N_{\text{TP}}N_{\text{PP}})$, enabling training of models that exceed single-device memory. The cost is extra communication for activations and gradients plus idle time from sequential dependencies, so the memory gain must exceed the transfer and pipeline-bubble overhead introduced by the partition.
2. **Distinction**: Unlike data parallelism, which replicates the full model on every worker and shards the minibatch, model parallelism shards the model itself. Pipeline parallelism and tensor parallelism are implementation forms of model parallelism: one partitions layers into stages, while the other partitions individual tensor operations.
3. **Common pitfall**: A frequent misconception is that model parallelism automatically speeds up training. Its possible gains are capacity, less traffic on the slow tier, and shorter layer time, and whichever gain a plan relies on must justify the added transfer and idle cost. Without careful scheduling and overlap, downstream devices wait for upstream activations and utilization can fall below a single-device baseline.

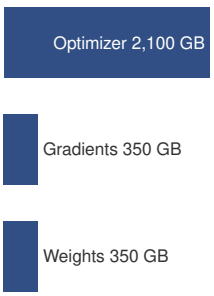

Weights are only a small slice of the full Adam training state.

Memory fit alone does not settle whether a model needs this kind of partition. For a 175B-parameter model, weights alone occupy 350 GB, or about 5.5 GB per GPU across a 64-GPU group, 1/16 of the 1,024-GPU reference cluster. Full mixed-precision Adam training state is much larger, 2,800 GB globally, yet ZeRO-3 across the same 64 GPUs reduces it to roughly 43.8 GB per GPU before activations, which fits an 80 GB accelerator. ZeRO-3 can therefore make the static state fit, as the notebook below confirms. It pays for that fit in traffic, because every step regathers every layer's parameters across the data-parallel fabric, roughly three times the model's parameter bytes per worker (section 6.3.3). It also cannot split an individual layer's computation or activations across devices.

Activations add a separate memory requirement. Under the earlier unchecked-storage assumptions, a 2048-token sequence through 175B parameters produces roughly 1.1 TB at a local batch of 4 sequences; optimizer sharding does not reduce that storage. Recomputation trades memory for extra FLOPs, partitioning distributes activations, and offload moves them to another memory tier. Kernel and batch-size choices also change the budget. Full activation checkpointing brings the same modeled case down to ~19.3 GB. The following notebook, which uses a single-sequence local batch, checks whether ZeRO-3 plus checkpointing fits the model without partitioning its computation.

Napkin Math 6.7: The memory wall of scale

Problem: Training a 175B-parameter model (like GPT-3) on a 64-GPU group of NVIDIA A100s (80 GB), smaller than the 1,024-GPU reference cluster. The A100 carries the same 80 GB of HBM as the reference H100, so the capacity verdict carries over, and a small group is the harder test of whether sharding alone suffices. Can data parallelism with ZeRO-3 handle this?

Math:

1. **Parameter storage**: 175B params $\times$ 2 bytes (FP16) = 350 GB.
2. **Gradients and optimizer state**: Gradients add 175B params $\times$ 2 bytes = 350 GB, while Adam FP32 master weights, momentum, and variance add 175B params $\times$ 12 bytes = 2,100 GB.
3. **Total static memory**: 2,800 GB.
4. **ZeRO-3 sharding**: With 64 GPUs, per-GPU static memory = 2,800 GB/64 GPUs $\approx$ 43.8 GB.

5. **Activation memory**: For sequence length 2048 and batch size 1, the canonical full-checkpointing strategy stores $\approx$ 4.8 GB of activations per GPU.

Systems insight: 43.8 GB (static) + 4.8 GB (dynamic) = 48.6 GB. Under this simplified accounting, ZeRO-3 plus full activation checkpointing fits within the 80 GB A100 capacity, leaving material headroom for framework buffers and transient layer gathers. If those additional allocations exceed the remaining HBM, the design can reduce the local batch, offload state, increase the sharding degree, or add tensor or pipeline parallelism; tensor parallelism is not the only recourse.

In this example, sharding and checkpointing make the modeled state fit, leaving two distinct reasons to partition the model itself. One is regather traffic. On the 1,024-GPU reference cluster, ZeRO-3 moves about 1.1 TB per worker per step for this model, which on one NDR InfiniBand port takes about 3× as each GPU's computation for one training step, and longer still when gradient accumulation regathers the parameters for every microbatch. The other is layer execution, since full sharding gathers parameters on demand without ever splitting a single layer's computation or activations across devices.

Model parallelism instead cuts the computation graph directly, either between sequential layer groups (such as layers 1–4 and layers 5–8) or inside a layer across channels and attention heads. Layer cuts introduce pipeline dependencies that leave downstream devices idle while an upstream stage computes, creating the pipeline bubble that governs utilization for this partition (section 6.5.2). Channel and head cuts, the tensor parallelism of section 6.5.3, address both regather traffic and layer execution by partitioning individual weight matrices across devices. They expose concurrent work within a layer at the price of intra-layer exchanges, which the column-then-row pairing of figure 6.16 confines to block boundaries.

These cuts enable the large-scale cases that exceed a single device for different reasons. GPT-3, with 175B parameters, was trained with model parallelism within each matrix multiplication and across layers (Brown et al. 2020). That partition avoids the full-parameter regather that ZeRO-3 pays on every step. Vision transformers processing high-resolution $16k \times 16k$ pixel images use model parallelism to bound activation memory. Mixture-of-experts (MoE) architectures, whose sparsely activated feed-forward blocks are developed in section 6.5.5, distribute their conditional computation paths across hardware (Shazeer et al. 2017; Lepikhin et al. 2021; Fedus et al. 2022).

Each partitioning strategy trades single-device memory overflow for a different systems cost: layer-wise partitioning pays in idle time, pipeline parallelism pays in scheduling complexity to hide that idle time, and tensor parallelism pays in NVLink bandwidth to split individual layers. Sequence and context parallelism, expert parallelism, and sparse embedding sharding then extend the same trade-off to long-context activations, sparsely activated experts, and sparsely touched embedding tables.

Checkpoint 6.3: Model parallelism foundations

Verify your understanding of model sharding:

- ☐ **Memory relief**: determine whether model parallelism reduces the memory footprint per GPU for a given model.
- ☐ **Sequential dependency**: identify whether sequential dependencies between layer-wise model shards arise in the forward pass, the backward pass, or both.
- ☐ **Pipeline bubble**: explain why unscheduled layer-wise model parallelism can lead to lower GPU utilization than data parallelism through the "pipeline bubble."
- ☐ **Layer vs. tensor split**: distinguish layer-based and tensor-based ways to split a Transformer layer.

6.5.1 Layer-wise partitioning

Where ZeRO-3 met the capacity limit by sharding storage while every worker still executed every layer, layer-wise model parallelism gives each device a *unique* contiguous block of layers to compute. It is the least invasive model cut: with 8-way partitioning, the 175B-parameter model's 350 GB of FP16 weights fall to approximately 44 GB per GPU, though activations, gradients, and optimizer state still need sharding, offload, or pipeline scheduling (section 6.6).

Because consecutive layer blocks remain sequentially dependent, however, a single batch passing through those stages in order leaves every other GPU waiting for its turn. In a 16-layer transformer placed four consecutive layers per GPU (figure 6.13), panel (a) shows each device holding one-quarter of the weights and exchanging only boundary activations and gradients, while panel (b) reveals the utilization penalty: without microbatch pipelining, each GPU computes in only two of the eight schematic time slots and sits idle for the other six.

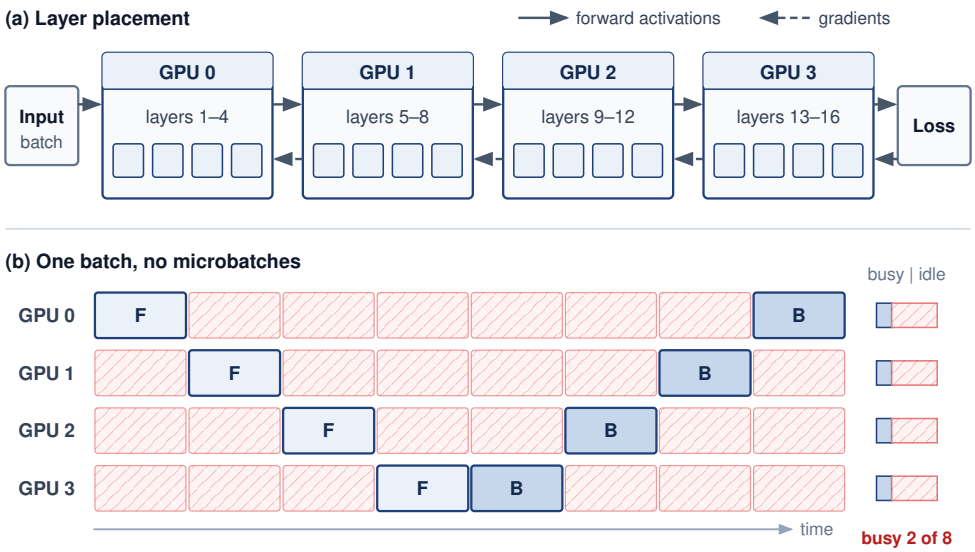

Figure 6.13: Layer-Wise Model Parallelism: (a) A 16-layer transformer placed four consecutive layers per GPU, cutting each GPU's weight footprint to one-quarter of the model. Forward activations cross each stage boundary left to right; gradients return right to left. (b) Without microbatching, one batch visits the stages in turn, forward (F) and then backward (B) in reverse, in equal schematic time slots, so each GPU works in only 2 of 8 slots. The microbatched schedule in figure 6.14 fills much of that idle time.

Partitioning layers changes where an optimizer update runs without necessarily changing the update rule. Each GPU can apply Adam (Kingma and Ba 2015), AdaFactor, or another optimizer to its parameter slice once the required gradients are complete. A device holding transformer layers 1–12 updates those layers. Tied parameters, which share one weight matrix across stages, can still require coordination, as can global gradient clipping, whose norm spans every stage, and a combined data-parallel group. Preserving synchronous optimization also requires the schedule to use consistent parameter versions. Sequential dependencies introduce pipeline bubbles, and unequal stage service times add waiting. The engineering problem is to overlap microbatches and balance stage work while preserving that update contract (Rasley et al. 2020).

6.5.2 Pipeline parallelism

Without pipelining, a four-stage layer partition leaves three of the four GPUs idle at any moment because downstream stages cannot start until upstream stages finish the entire batch (figure 6.13). Pipeline parallelism breaks that stall by splitting the batch into smaller microbatches and streaming them through the stages like a relay race. As soon as GPU 1 finishes the initial transformer blocks on microbatch 1, it hands the resulting activation tensor to GPU 2 and immediately starts microbatch

2. For the 175B model with hidden dimension 12,288 and a microbatch of 4 sequences at 2,048 tokens each, that forward handoff carries approximately 201 MB per stage boundary per microbatch. GPU 2 needs that payload before processing the same microbatch, creating a dependency chain through the stages, while other microbatches occupy the remaining stages concurrently. Backward propagates activation gradients in reverse as each stage computes gradients for its own parameters.

Definition 6.8: Pipeline parallelism

Pipeline parallelism is a model parallelism technique that partitions a neural network's layers into sequential stages assigned to different devices, passing activations forward and gradients backward between stages while overlapping computation across stages using microbatches to maintain throughput.

1. **Significance**: Each boundary transfers forward activations and backward activation gradients. A forward tensor contains $B_{\text{micro}} \times S \times d \times 2$ bytes at BF16, where B_{micro} is the microbatch size. For a hidden dimension of 8,192 with microbatch size 1 and a 2,048-token sequence, the forward payload is approximately $8{,}192 \times 2{,}048 \times 2 \approx 34$ MB per microbatch per boundary. Backward traffic and all microbatches must be counted before comparing this with a per-update gradient AllReduce or sizing a 50 GB/s InfiniBand link. In a balanced noninterleaved schedule with ideal transfers, the idle fraction is approximately $(N_{\text{PP}} - 1)/(m + N_{\text{PP}} - 1)$, where N_{PP} is the stage count and m the microbatch count.
2. **Distinction**: Tensor parallelism partitions operations within a layer and introduces communication among their participants. Pipeline parallelism instead transfers activations and their gradients between sequential stages. This coarser exchange can suit slower links when the transfer fits within the stage schedule; the interconnect name alone does not establish feasibility.
3. **Common pitfall**: A frequent misconception is that adding pipeline stages always improves throughput. Each additional stage increases the pipeline bubble fraction $(N_{\text{PP}} - 1)/(m + N_{\text{PP}} - 1)$ unless the number of microbatches m scales proportionally, causing hardware to idle and erasing the benefit of deeper partitioning.

In the GPipe schedule of figure 6.14, a four-stage pipeline processes four microbatches ($N_{\text{PP}} = 4, m = 4$). Dividing the global batch into discrete microbatches allows forward passes to overlap diagonally across stages before backward passes unfold in reverse order. While microbatching raises hardware utilization above naive sequential placement, pipeline fill (t_1–t_3) and drain (t_{12}–t_{14}) bubbles still leave each device idle for 6 of 14 time slots.

Tracing the schedule grid in figure 6.14 column by column reveals how the fill and drain bubbles arise. In the first slot only stage 0 has work, the forward pass of microbatch 1. In the second, stage 0 starts microbatch 2 while stage 1 receives microbatch 1's activations, and by the fourth slot every stage holds a different microbatch. The last stage then turns microbatch 4 around into its backward pass, and gradients travel upstream in reverse microbatch order, so stage 0 waits until the eleventh slot for its first gradient. Each stage therefore idles 3 slots in the forward phase and 3 more in the backward phase, 6 of the 14 slots in the step. Those idle slots may fall at the start of its row, the middle, or the end. That share, about 43 percent, is the bubble fraction $(N_{\text{PP}} - 1)/(m + N_{\text{PP}} - 1)$ at $N_{\text{PP}} = m = 4$. Adding microbatches lengthens the busy diagonal without lengthening fill and drain in this balanced schedule. Increasing m is therefore one way to reduce the idle fraction; later schedules also change which work can occupy the gaps.

The same bubble term becomes concrete in the following worked example, where an 8-stage pipeline and 32 microbatches still leave a measurable amount of idle time.

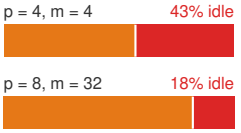

Enough microbatches shrink the bubble even as stages double.

↪ The bubble fraction across stage and microbatch counts, and its interleaved-schedule variant, are developed in section D.4.

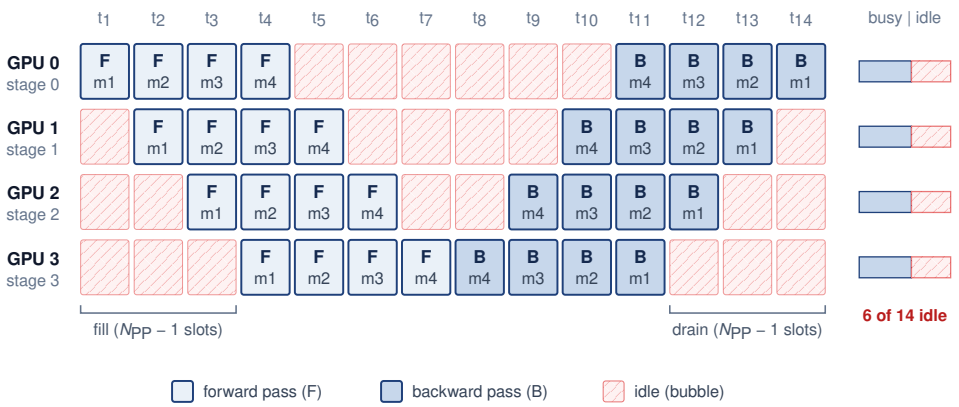

Figure 6.14: Pipeline Parallelism Schedule: A 4-stage pipeline processes 4 microbatches with all forward passes (F) before any backward pass (B), as in GPipe. Rows are pipeline stages and columns are equal time slots. Each backward pass enters a stage only after the next stage has produced that microbatch's gradient, so every stage idles $N_{\text{PP}} - 1$ slots in the forward phase and $N_{\text{PP}} - 1$ in the backward phase, 6 of 14 slots in all.

Napkin Math 6.8: The cost of the pipeline bubble

Problem: A large-model training run uses pipeline parallelism across 8 stages. To hide the sequential delay, the batch is split into 32 microbatches. What fraction of GPU cycles does the pipeline bubble leave idle?

Math: For balanced, noninterleaved synchronous stages under the one-forward-one-backward (1F1B) schedule developed below, with ideal transfers, the bubble fraction follows from the stage and microbatch counts. The time units below each combine one forward and one backward service interval; they are not individual optimizer steps.

1. **Wait time**: At the start and end of each batch, GPUs sit idle for $N_{\text{PP}} - 1$ steps.
2. **Productive time**: GPUs compute for m steps.
3. **Bubble fraction**: $(N_{\text{PP}} - 1)/(N_{\text{PP}} - 1 + m) = 7 \,/\, (7 + 32) \approx 17.9$ percent.

Systems insight: Pipeline parallelism is a *utilization-memory trade-off*. With this stage count and schedule fixed, reducing the bubble from 18 percent to 5 percent requires more microbatches (m). In general, holding the bubble under 10 percent requires $m > 9(N_{\text{PP}} - 1)$, because $(N_{\text{PP}} - 1)/(m + N_{\text{PP}} - 1) < 0.1$ rearranges to that bound. The activation memory those microbatches need depends on the schedule. GPipe retains activations across the microbatches, whereas the 1F1B schedule developed next releases them during the steady state. The modeled 17.9 percent idle fraction also motivates interleaved pipelining, which divides each device's work into smaller stages. Interleaving can reduce the original 18 percent bubble, but additional transfers and memory requirements determine the net gain.

GPipe[4] (Huang et al. 2019) introduced synchronous pipeline parallelism with microbatch accumulation, while PipeDream (Narayanan et al. 2019) developed asynchronous approaches with weight stashing, which keeps one parameter version per in-flight minibatch so that its forward and backward passes use the same weights. Modern systems use 1F1B (one-forward-one-backward)[5] scheduling to reclaim activation memory earlier. That lower memory pressure makes larger microbatch counts practical, and the larger m is what shrinks the bubble term.

At the schedule level, 1F1B is a warmup, steady-state, and drain loop. In algorithm 6.1, each stage alternates forward work for newer microbatches with backward work for older microbatches once gradients arrive, so activations can be released before the whole batch has completed.

[4] **GPipe**: Published by Google in 2019, GPipe introduced synchronous microbatch pipelining that trained a 557M-parameter AmoebaNet across 8 Tensor Processing Units (TPUs) with near-linear scaling. Its schedule keeps every microbatch's activations live until the backward pass, the cost that later schedules reduce.

[5] **1F1B (one-forward-one-backward) scheduling**: Unlike GPipe (which processes all forward passes before any backward passes), 1F1B interleaves them. For equal-sized microbatch activations and a non-interleaved schedule, peak retained activations per stage fall from $\mathcal{O}(m)$ under GPipe to at most $\mathcal{O}(N_{\text{PP}})$ under 1F1B, where m is the microbatch count and N_{PP} the pipeline-stage count. Summed across stages, the corresponding bounds are $\mathcal{O}(mN_{\text{PP}})$ and $\mathcal{O}(N_{\text{PP}}^2)$. This memory reclamation is what enables the large microbatch counts $(m \gg N_{\text{PP}})$ required to keep the pipeline bubble small.

Algorithm 6.1 One-forward-one-backward pipeline schedule

Require: N_{PP} ordered pipeline stages; m microbatches; activation links $j \to j{+}1$ and gradient links $j{+}1 \to j$
Ensure: one optimizer update after gradients accumulate across all m microbatches
1: split the global batch into m microbatches; assign consecutive layers to the N_{PP} stages
2: warm-up: launch forwards from stage 0, forwarding each activation downstream as produced
3: **for** each stage in steady state **do**
4: alternate one backward (oldest ready microbatch) with one forward (next microbatch)
5: release a microbatch's activation as soon as that stage's backward consumes it
6: **end for**
7: drain: finish the remaining backward passes in reverse stage order
8: apply the optimizer update once every stage has accumulated gradients for all m microbatches

Run on four stages with eight microbatches, the schedule gives each stage a short warm-up of forward passes, a steady state that alternates one backward with one forward, and a drain of the remaining backward passes (figure 6.15). The warm-up and drain phases still leave a bubble of roughly $(N_{\text{PP}}-1)/(m+N_{\text{PP}}-1)$, 6 of 22 slots here. Interleaving backward work in the steady state, however, frees each activation as soon as its backward consumes it. Stage 0 never holds more than four microbatches' activations where an all-forwards-first schedule would hold all eight, so peak activation memory grows with the stage count rather than the microbatch count. That is the gain over GPipe-style all-forward-then-all-backward scheduling, and it lets the system raise m until activation memory, recomputation, or the global-batch constraint becomes the next limit.

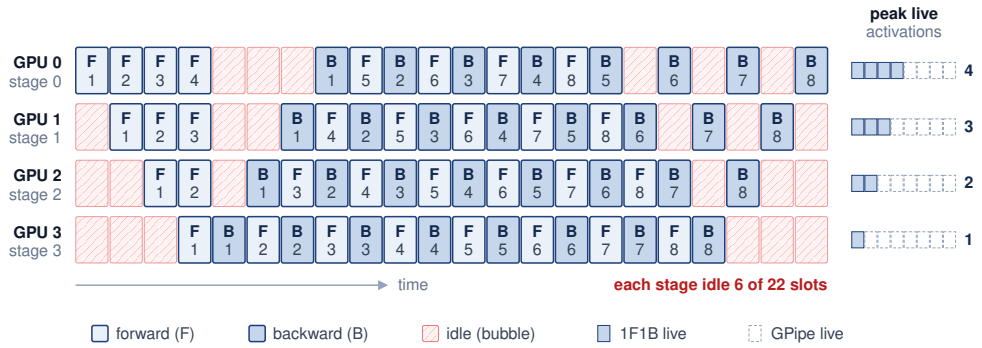

Figure 6.15: One-Forward-One-Backward (1F1B) Pipeline Schedule: Four stages process eight microbatches in equal schematic slots. After a warm-up of forward passes, each stage alternates one backward pass with one forward pass, then drains its remaining backward passes. Fill and drain still idle every stage for 6 of 22 slots, the bubble fraction $(N_{\text{PP}}-1)/(m+N_{\text{PP}}-1)$, but each stage releases a microbatch's activations as soon as that microbatch's backward pass runs. The bars at the right count the most microbatches whose activations a stage holds at once, four at stage 0 and one at the last stage, against the eight an all-forwards-first schedule would hold (dashed).

The activation handoff has shape $B_{\text{micro}} \times S \times d$; its completion still constrains the next stage. The 1F1B schedule in algorithm 6.1 reduces activation residency, but fill, drain, communication, and load imbalance leave idle time.

Zero-bubble pipeline schedules further reduce idle time by decoupling the two matrix multiplications inside each layer's backward pass. In a standard backward pass, the GPU computes both $\partial\mathcal{L}/\partial\mathbf{X}$ (the activation gradient propagated backward to the preceding stage) and $\partial\mathcal{L}/\partial\mathbf{W}$ (the weight gradient accumulated locally for the optimizer step) back-to-back as a single backward unit. Zero-bubble scheduling splits them into separate kernels. A B kernel multiplies the incoming upstream gradient by the stage's persistent weights $\mathbf{W}^T$ to compute only the activation gradient $\partial\mathcal{L}/\partial\mathbf{X}$ and immediately sends it to the previous stage. A W kernel multiplies the saved forward activation $\mathbf{X}^T$ by the upstream gradient to compute the weight gradient $\partial\mathcal{L}/\partial\mathbf{W}$ locally. Because upstream stages cannot begin their own backward work until they receive $\partial\mathcal{L}/\partial\mathbf{X}$, the B kernel sits

on the critical path and must execute promptly, whereas the W kernel has no cross-stage dependency and can be scheduled opportunistically to fill local pipeline bubbles.

The B/W split creates scheduling freedom, although all deferred W kernels must still finish before the corresponding optimizer update. In the balanced four-stage, eight-microbatch schedule of figure 6.15, noninterleaved 1F1B still leaves $3/11 \approx 27$ percent of each stage's time idle, by the bubble fraction derived above. Reordering forward, activation-gradient, and weight-gradient work allows deferred W kernels to soak up those idle slots. Approaching zero bubble time in practice also depends on kernel durations, memory headroom, and optimizer synchronization, so splitting the backward pass alone does not guarantee it.

The trade-off is memory. Because the B kernel needs only the incoming gradient and the persistent weights **W**, whereas the W kernel requires the saved forward activation **X**, a stage cannot free a microbatch's forward activation when B finishes. Deferring W to fill a later bubble forces the stage to keep both **X** and the incoming gradient resident in device memory until W finally executes, increasing peak memory usage. Some implementations address this by combining zero-bubble scheduling with activation checkpointing, selectively recomputing certain activations before W rather than storing them across the delay. The interaction between these techniques creates a three-way trade-off among pipeline bubble size, memory consumption, and recomputation overhead, an example of the displacement of overhead (principle 14).

6.5.3 Tensor parallelism

Every pipeline schedule so far keeps whole layers on each device, which is why its traffic is confined to stage boundaries and tolerates moderate interconnect bandwidth. The same property sets its limit. A pipeline cannot help when a single layer's weights or activations exceed one device, and it cannot shorten the time one layer takes, because that layer still runs on one accelerator.

Tensor parallelism splits the weight matrices within each layer instead of assigning complete layers to devices. This operator-level parallelism (also called intra-layer parallelism) enables finer-grained distribution but requires high-bandwidth interconnects for the frequent intra-layer communication it introduces.

Definition 6.9: Tensor parallelism

Tensor parallelism partitions individual tensor operations across devices, commonly through column-parallel or row-parallel weight splits. Communication combines or redistributes the resulting tensor shards as required by the next operation. The Megatron-style forward-pass example below illustrates one such arrangement.

1. **Significance**: The Megatron-LM arrangement developed below uses two AllReduce operations per transformer block in the forward pass, one after attention and one after the MLP. Each reduces a 2-byte activation tensor of $B \times S \times d$ elements, where B counts the sequences in flight through the layer. A ring AllReduce moves roughly $2(N_{\text{TP}} - 1)/N_{\text{TP}}$ times that payload per GPU, and the exchange recurs in every one of the 96 blocks of the GPT-3 shape, with backward adding corresponding exchanges. Whether that traffic hides behind computation depends on the link that carries it, which the arithmetic after the mechanism quantifies.
2. **Distinction**: Unlike pipeline parallelism, which communicates only at layer boundaries and transfers small activation slices, tensor parallelism synchronizes within every layer via AllReduce. It is therefore sensitive to the per-operation latency of the interconnect and favors a fast communication domain such as NVLink, which can extend beyond one node.
3. **Common pitfall**: Extending tensor parallelism across nodes without checking the transfer budget can erase the compute saving. The reference NVLink specification is approximately 900 GB/s bidirectional, or 450 GB/s per direction; InfiniBand NDR provides 50 GB/s per port, a 9× per-direction gap. These specifications identify a placement risk. Whether the split helps depends on sustained bandwidth, operation size, compute saved, and overlap.

Frequent intra-layer exchanges favor placing tensor-parallel groups within the fastest available communication domain. Pipeline stages communicate at coarser boundaries and can often span slower domains when their transfers fit the schedule. This difference motivates a hybrid that maps each exchange to a suitable fabric tier; it does not impose a universal node boundary for either strategy.

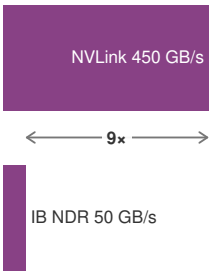

Tensor-parallel exchanges belong on the fast tier, about 9× an NDR port.

Megatron-style tensor parallelism[6] (Shoeybi et al. 2019) partitions matrix multiplications by alternating column and row splits. In the two-GPU dataflow of figure 6.16, column-parallel splitting partitions the first weight matrix $\mathbf{A} = [\mathbf{A}_1 \mid \mathbf{A}_2]$ to produce local activation slices $\mathbf{Y}_1, \mathbf{Y}_2$ from replicated input $\mathbf{X}$. Row-parallel splitting then partitions the second weight matrix $\mathbf{B} = [\mathbf{B}_1; \mathbf{B}_2]$ so that each GPU computes partial products $\mathbf{Z}_1, \mathbf{Z}_2$ locally. This defers cross-device communication to a single AllReduce that reconstructs output $\mathbf{Z}$.

Column-parallel linear layers split weights along columns. For input $\mathbf{X}$ and weight matrix $\mathbf{W} = [\mathbf{W}_1 \mid \mathbf{W}_2]$ split across 2 GPUs:

$$\mathbf{Y} = \mathbf{X}\mathbf{W} = \mathbf{X}[\mathbf{W}_1 \mid \mathbf{W}_2] = [\mathbf{X}\mathbf{W}_1 \mid \mathbf{X}\mathbf{W}_2]$$

Each GPU computes its partition independently. The concatenation in the equation is logical, because when the next operation is row-parallel, each output shard remains on its producing device and feeds the matching weight shard without an AllGather. The row-parallel half of the pattern splits weights along rows, $\mathbf{W} = \begin{bmatrix} \mathbf{W}_1 \\ \mathbf{W}_2 \end{bmatrix}$, and the layer's input arrives split along its columns to match, each GPU holding the shard it kept from a preceding column-parallel layer:

$$\mathbf{Y} = \mathbf{X}\mathbf{W} = \mathbf{X}_1\mathbf{W}_1 + \mathbf{X}_2\mathbf{W}_2$$

Each GPU computes a partial sum, and the system pays one AllReduce to combine the partial outputs.

6 **Megatron-LM**: NVIDIA's 2019 framework that trained an 8.3B-parameter transformer (24× BERT-Large and 5.5× GPT-2 at the time) with only two AllReduce operations per transformer block in the forward pass (one after attention, one after the multilayer perceptron (MLP)). The run reached 76 percent scaling efficiency across 512 GPUs, and its partitioning remains influential in large-scale transformer training.

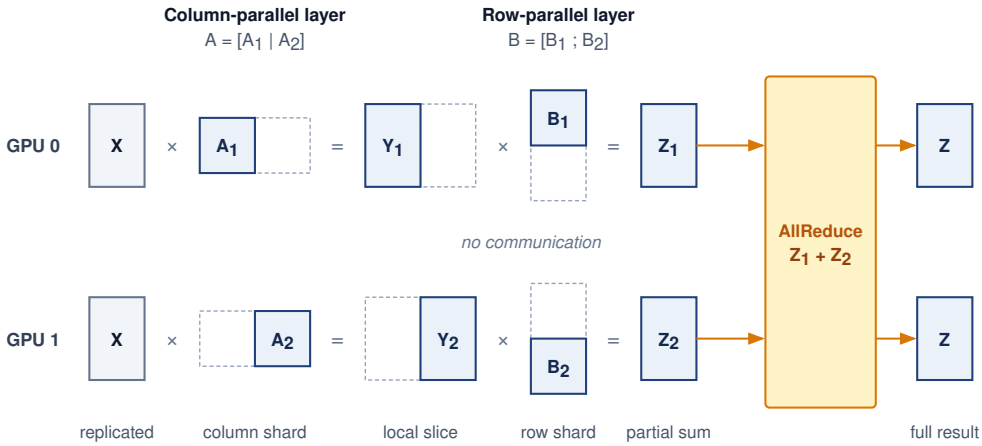

Figure 6.16: Matrix Partitioning under Tensor Parallelism: A Megatron-LM layer pair on two GPUs, read left to right. The first layer (for example, QKV or the first feed-forward layer) is split by columns, $\mathbf{A} = [\mathbf{A}_1 \mid \mathbf{A}_2]$, so each GPU computes its own slice $\mathbf{Y}_i$ of the replicated input's output. The second layer is split by rows, $\mathbf{B} = [\mathbf{B}_1; \mathbf{B}_2]$, so each slice feeds its matching shard on the same GPU without communication. The only collective is the amber AllReduce that sums the partial outputs $\mathbf{Z}_1 + \mathbf{Z}_2$, leaving the full result on both GPUs.

The column-then-row sequence of figure 6.16 embodies the governing architectural invariant. Pairing a column-parallel layer directly with a row-parallel layer allows intermediate activations $(\mathbf{Y}_1, \mathbf{Y}_2)$ to remain strictly local on each GPU, confining inter-device synchronization to a single collective at the end of the block. This locality holds because the intermediate operation between the two projections acts independently on each GPU's column slice without cross-rank data.

Megatron applies that pattern twice in each transformer block. In multi-head attention, the QKV projection is column-parallel along the head dimension so that each GPU receives the complete query, key, and value vectors for a disjoint subset of $N_{\text{heads}}/N_{\text{TP}}$ attention heads, computes attention locally for those heads, and feeds the resulting head outputs into the row-parallel attention output projection. In the feed-forward network, the first linear layer is column-parallel, the element-wise nonlinearity applies independently to each coordinate of the local slice, and the second linear layer is row-parallel. The forward pass therefore places one AllReduce after attention and one after the feed-forward network, for two forward AllReduce operations per transformer layer. In the backward pass, the dataflow reverses so that gradients enter each row-parallel output layer without communication, while at the input of each column-parallel layer every GPU computes a partial input gradient from its local weight slice and runs an AllReduce to sum those contributions before the preceding residual connection, adding two symmetric backward AllReduces per layer.

Communication volume per transformer layer depends on sequence length S, hidden dimension d, tensor-parallel degree N_{TP}, and batch size B. The batch size is the number of sequences in flight through the layer (the local batch, or the microbatch B_{micro} when the layer also sits in a pipeline). The activation payload is:

$$M_{\text{act}} = B \times S \times d \times s_{\text{elem}}$$

A ring AllReduce transfers approximately $2(N_{\text{TP}}-1)/N_{\text{TP}} \times M_{\text{act}}$ bytes per GPU, which is close to $2M_{\text{act}}$ for large N_{TP}. For the GPT-3 shape ($S = 2048$, $d = 12{,}288$, 96 layers) with $B = 4$ and FP16, the activation payload is $4 \times 2048 \times 12{,}288 \times 2 \approx 201$ MB.

At $N_{\text{TP}} = 8$ the ring factor is 1.75, giving roughly 352 MB per AllReduce. With two AllReduces in each of the 96 layers, this totals approximately 67.6 GB per forward microbatch. The symmetric backward exchanges bring the total to roughly 135 GB per forward/backward microbatch, before any recomputation. Set against the computation it serves, that volume is modest on the right link and prohibitive on the wrong one. At the per-layer compute rate of the earlier feasibility check, one such microbatch takes about 2.2 s of forward and backward computation per GPU. Its tensor-parallel traffic takes about 0.3 s on NVLink but about 2.7 s on one InfiniBand port, longer than the computation. The gap matters because this traffic sits on the critical path of every layer. The next operation consumes each reduced result and cannot start on that data until the reduction completes, whereas a data-parallel gradient reduction can overlap the backward computation that remains. The number of microbatches per update and the available overlap then determine the required fabric service rate. Computing MFU requires the complete compute and exposed-communication time, not the bandwidth ratio alone.

Tensor parallelism eventually reaches a workload-dependent limit because smaller operation shards leave less computation to cover communication. Per-GPU ring AllReduce bytes approach $2\times$ the activation tensor size, while local matrix shapes and collective startup costs change with group size. The useful degree therefore depends on layer shape, microbatch size, and the fast communication domain. Llama 3 405B's TP=8 arrangement within NVLink-connected H100 nodes (Dubey et al. 2024) and node-local groups in other large training systems (Z. Jiang et al. 2024) illustrate a placement choice, not a universal eight-device limit. Tensor parallelism also costs more to adopt than data parallelism. It replaces standard linear layers with column-parallel and row-parallel variants inside the model definition, whereas data parallelism wraps the model externally without modifying its internals.

6.5.4 Sequence and context parallelism

While tensor parallelism partitions weight matrices within linear layers, the memory footprint of large-scale transformer training is increasingly dominated by activations. As models scale their context length from thousands to hundreds of thousands or millions of tokens ($S \geq 32\text{k} \ldots 1\text{M}+$), the activation footprint scales proportionally:

$$M_{\text{activations}} = \mathcal{O}(B \times S \times d \times N_L)$$

for linear, normalization, and residual layers, and $\mathcal{O}(B \times N_{\text{heads}} \times S^2 \times N_L)$ for un-tiled attention (or $\mathcal{O}(B \times S \times d \times N_L)$ with tiled FlashAttention, which computes attention block by block without storing the full score matrix). At long contexts, activation memory dwarfs parameter storage, exposing the memory capacity gap (principle 2). Two complementary methods address this constraint by

sharding along the sequence axis: sequence parallelism (SP) within the tensor-parallel domain, and context parallelism (CP) across dedicated long-sequence communication groups.

Definition 6.10: Sequence and context parallelism

Sequence and context parallelism partition the token sequence dimension to reduce per-device activation storage. Sequence parallelism (SP) retains sequence shards during operations such as LayerNorm and Dropout within a tensor-parallel group. Context parallelism (CP) partitions the sequence across a dedicated group and exchanges the information needed for attention, through mechanisms developed below that either pass key/value blocks around a ring or transpose between sequence and head partitions.

1. **Significance**: Long contexts can exhaust activation memory even when sharded model state fits. The relevant limit depends on batch size, layer shapes, attention kernels, and checkpointing. Sequence parallelism divides the targeted activation storage by N_{TP} while preserving the ideal collective payload volume in the decomposition below. Context parallelism distributes sequence-dependent storage; linear context growth ($c \times S_{\text{single}}$, with c denoting the context-parallel degree) is an ideal capacity estimate that still requires communication buffers and computation to fit the budget.
2. **Distinction**: Tensor parallelism partitions weight matrices along hidden dimensions d, whereas context parallelism partitions the sequence axis S. Ring Attention exchanges key/value chunks between devices, while Ulysses uses an AllToAll to transpose between sequence and head dimensions. Both require placement that can service their exchanges within the attention schedule.
3. **Common pitfall**: Assuming context parallelism eliminates the quadratic attention compute complexity. CP distributes the attention computation across devices, reducing *per-device* memory and FLOPs by $1/c$, but the *aggregate* cluster FLOPs for dense self-attention still scale quadratically as $\mathcal{O}(S^2)$.

6.5.4.1 Megatron sequence parallelism

In the Megatron-style arrangement just derived, the column-then-row pairing in figure 6.16 splits the QKV projection and first MLP layer column-wise, and the attention output projection and second MLP layer row-wise. The surrounding LayerNorm, Dropout, and residual operations retain replicated activations across all N_{TP} devices. LayerNorm operates across a token's hidden features, while Dropout and residual addition operate element-wise; none requires another token's features. This token independence permits partitioning the otherwise replicated $B \times S \times d$ activation tensor along its sequence dimension.

Megatron Sequence Parallelism (Korthikanti et al. 2023) eliminates this redundancy by sharding the sequence dimension across the N_{TP} ranks for all non-tensor-parallel operations. Instead of replicating activations, each GPU holds only a $B \times (S/N_{\text{TP}}) \times d$ slice during LayerNorm and Dropout.

This sharding incurs zero additional communication volume. In standard tensor parallelism, the row-parallel matrix multiplication is followed by an AllReduce to sum partial products across ranks. An AllReduce decomposes into a ReduceScatter followed by an AllGather, the identity established among the collective primitives in section 5.4.1:

$$\text{AllReduce}(x) \equiv \text{AllGather}(\text{ReduceScatter}(x))$$

Megatron-SP decouples this collective pair across the transformer block. Following the row-parallel linear layer, a ReduceScatter reduces partial sums while scattering the resulting activations along the sequence dimension, leaving each GPU with an S/N_{TP} token slice. LayerNorm, residual additions, and Dropout then execute on these local slices with zero cross-device communication. Immediately before the next column-parallel linear layer, an AllGather across the sequence dimension reassembles the complete S tokens required for the weight-partitioned matrix multiplication.

In this ideal payload accounting, standard TP executes two AllReduces per transformer block ($2 \times 2(N_{\text{TP}} - 1)/N_{\text{TP}} \times M_{\text{act}}$ bytes), while Megatron-SP executes two ReduceScatter and AllGather

pairs ($2\times[(N_{\text{TP}}-1)/N_{\text{TP}}+(N_{\text{TP}}-1)/N_{\text{TP}}]\times M_{\text{act}}=2\times 2(N_{\text{TP}}-1)/N_{\text{TP}}\times M_{\text{act}}$ bytes). The volumes match while the targeted activation storage falls by a factor of N_{TP}. Equal bytes do not imply equal elapsed time, because separating the collectives changes readiness, startup costs, and overlap opportunities.

6.5.4.2 Context parallelism with ring attention versus Ulysses

Megatron-SP shards the S dimension across the N_{TP} devices only outside the attention and MLP blocks; inside self-attention, each tensor-parallel rank still processes the full sequence length S for its local heads, so long-context runs can still exhaust HBM even after sequence parallelism and activation checkpointing. Context parallelism of degree c breaks that ceiling by partitioning the sequence itself across c accelerators throughout the entire block, including attention. Because every query token must still attend to keys and values from other sequence partitions, context parallelism requires cross-partition exchanges, supplied either by ring-based point-to-point transfers, by AllToAll head-sequence transpositions, or by a hierarchical combination of both.

Ring Attention (Liu et al. 2023) organizes devices in a logical ring where each GPU holds a query chunk $\mathbf{Q}_i$ of length S/c and an initial key/value chunk $(\mathbf{K}_i,\mathbf{V}_i)$. Computation proceeds in c rounds along the ring. In each round, GPU i computes block attention between its stationary $\mathbf{Q}_i$ and resident $(\mathbf{K}_j,\mathbf{V}_j)$ tiles using an online softmax accumulator. At the same time, it asynchronously transmits $(\mathbf{K}_j,\mathbf{V}_j)$ to ring neighbor $(i+1)$ mod c and receives $(\mathbf{K}_{j-1},\mathbf{V}_{j-1})$ from $(i-1)$ mod c. After c rounds every query chunk has attended to every key/value tile, yet no device ever stores the full key/value sequence (figure 6.17). Normalizing a query's attention scores requires all of them, but they arrive one tile at a time. The online softmax therefore maintains a running maximum and running sum of attention scores as tiles arrive and rescales earlier partial sums whenever the maximum changes. Each round's partial result merges into that running state, which likewise spares each device the full attention matrix.

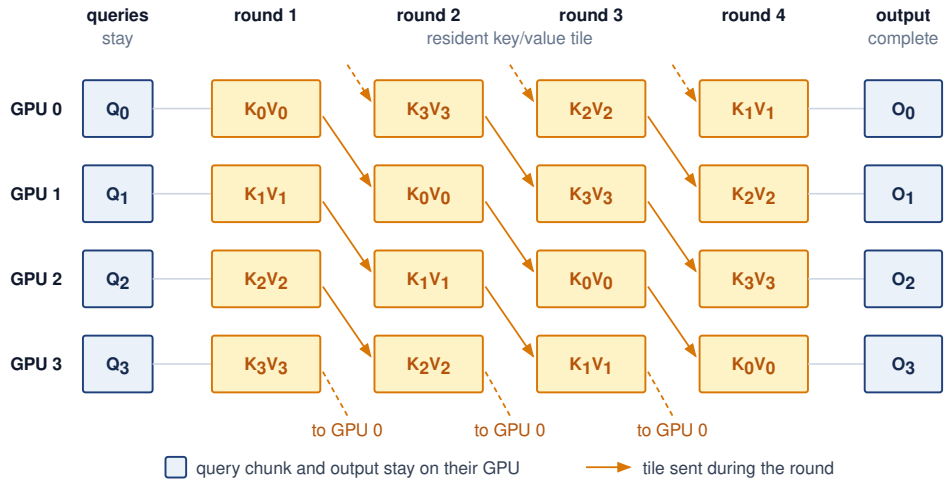

Figure 6.17: Ring Attention Rounds: Context parallelism of degree four on four GPUs. Each GPU keeps its query chunk and, in each round, attends to the key/value tile it currently holds while sending that tile to the next GPU in the ring; dashed arrows close the ring from GPU 3 back to GPU 0. After four rounds every query chunk has seen every tile, so each GPU's attention output for its own queries is complete. The transfer hides behind the block computation only when chunks are long enough, the overlap condition derived next.

For one sequence with standard multi-head attention, let $n=S/c$ be the tokens per chunk, d the total head width, s_{elem} the bytes per stored element, $R_{\text{peak}}\eta_{\text{hw}}$ the sustained attention FLOP/s, and BW_{link} the one-way transfer rate. Computing both $\mathbf{QK}^T$ and the weighted-value product costs

approximately $4n^2d$ FLOPs, while sending **K** and **V** costs $2nd\,s_{\text{elem}}$ bytes. Ignoring startup, softmax, and overlap interference gives the necessary tile-size estimate:

$$\frac{4n^2d}{R_{\text{peak}}\eta_{\text{hw}}} \geq \frac{2nd\,s_{\text{elem}}}{\text{BW}_{\text{link}}} \iff n \geq \frac{s_{\text{elem}}R_{\text{peak}}\eta_{\text{hw}}}{2\text{BW}_{\text{link}}}$$

This compares computation time with transfer time; HBM bandwidth cannot substitute for $R_{\text{peak}}\eta_{\text{hw}}$. It is an idealized overlap condition, not a guarantee.

For an H100 sustaining half its dense FP16 peak with 2-byte elements, the condition requires chunks of at least about 1,100 tokens per device when KV tiles travel over NVLink, and about 9,900 tokens when they cross one NDR InfiniBand port. A million-token context split across a few dozen devices clears both bounds, whereas splitting a 32K context across many nodes can leave each chunk too short to hide its transfers. Grouped-query attention, in which several query heads share one key-value head, changes the transferred KV width, and causal masking changes the tile work. Both require recomputing the two costs.

In causal autoregressive attention, a query token can attend only to earlier or current positions, so whenever a rank holding query chunk i receives a key and value chunk from a later sequence position $j > i$, the entire upper-triangular tile is masked out. Under naive contiguous partitioning, rank $i = 0$ holds the earliest tokens and performs useful work only on its own chunk ($j = 0$), sitting idle during the remaining $c-1$ ring rounds while rank $c-1$ (holding the final tokens) computes on every round. Because the ring steps synchronously at the pace of the busiest rank, this load imbalance leaves half the cluster's compute cycles idle. Production implementations eliminate this skew with zigzag or striped token partitioning, which splits the sequence into $2c$ sub-chunks and gives each device one early sub-chunk and one complementary late sub-chunk (for example, one chunk from the start of the sequence and one from the end). Every device then holds the same share of unmasked causal tiles in every round of the ring.

DeepSpeed Ulysses (Jacobs et al. 2023) exploits the independence of attention heads rather than circulating tiles along a ring. Each head attends over the whole sequence without reading any other head's scores, so the attention work can be split by heads instead of by tokens. Outside the attention core, activations are sharded along the sequence dimension (S/c tokens, all N_{heads} attention heads). Before computing self-attention, an AllToAll collective transposes the tensor from sequence-partitioned `[S/c, h, d_head]` to head-partitioned `[S, h/c, d_head]`. Each GPU now holds the full sequence length S, but only a subset N_{heads}/c of attention heads. Each device computes standard local FlashAttention across the full sequence for its heads; each rank owns complete heads, while the kernel handles causal masking locally. A second AllToAll collective transposes the attention output back to `[S/c, h, d_head]`. Ulysses is conceptually clean and reuses single-GPU attention kernels, but requires that the head count be divisible by the context degree and therefore no smaller than it ($N_{\text{heads}} \geq c$). Under grouped-query attention, where the key-value head count H_{KV} is smaller, a direct head-sharding implementation requires compatible divisibility and $c \leq H_{\text{KV}}$. Replicating KV heads changes that limit and its communication budget.

The ownership map in figure 6.18 makes this transpose explicit. Unlike the stationary query chunks in figure 6.17, a Ulysses rank receives every token for its assigned heads.

The exchanged tiles change where attention executes rather than removing its interactions. Head count and KV-head layout constrain this partition even when the activation bytes fit in memory.

A hierarchical configuration can combine head-sequence transposes within a fast domain and ring exchanges between domains. For an eight-device intra-node domain, a Ulysses group satisfies ($c_{\text{intra}} \leq 8$), while a second degree (c_{inter}) determines the ring's inter-node partition. In this example, Ulysses partitions heads locally and Ring Attention circulates KV blocks across nodes. The arrangement is feasible only if head divisibility, peak memory, link service time, and the overlap condition all hold; hierarchy alone does not prevent fabric saturation.

6.5.5 Expert parallelism (mixture of experts)

Tensor, pipeline, sequence, and context parallelism all partition dense layers that every token traverses, so scaling a dense model's parameter count increases its per-token compute cost in lockstep. MoE architectures (Shazeer et al. 2017) break that lockstep through conditional computation, and expert parallelism distributes the resulting capacity across devices. In an MoE layer, the feed-forward

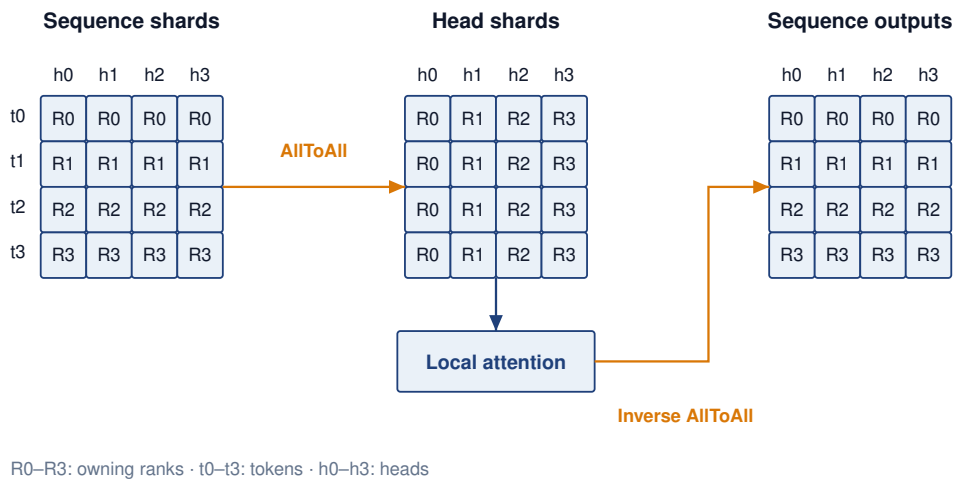

Figure 6.18: Ulysses Ownership Transpose: A schematic with four ranks, four tokens, and four attention heads labels each token–head tile by its owning rank. The first AllToAll changes ownership from token rows to head columns. Each rank computes attention across the full sequence for its head, then the inverse AllToAll restores token ownership for subsequent layers. The head dimension is omitted. This example assumes ordinary multihead attention with compatible divisibility; it does not depict KV-head replication under grouped-query attention.

block contains N_{experts} independent feed-forward networks ("experts"), and a gating network selects only a small subset, often top-1 or top-2, for each token. Adding more experts while holding each expert's shape and the selected count fixed expands stored parameters without increasing expert FLOPs per token, trading dense matrix work for routing and all-to-all communication.

In a distributed setting, experts are partitioned across workers, with each rank hosting a disjoint subset of expert networks. One routing cycle across four ranks shows the sequence (figure 6.19). After top-k_{experts} gating selects destination experts, an AllToAll dispatch shuffles tokens across the fabric to the hosting ranks. Expert FFNs then compute locally, and a second AllToAll return routes the processed activations back to the originating ranks to form the gate-weighted sum. The backward pass repeats both exchanges in reverse order, sending output gradients to the ranks hosting the experts and returning input gradients to the ranks that own the tokens. The primary advantage is decoupling model size from per-token compute, which lets parameter count grow within feasible hardware budgets. The primary constraint is the AllToAll communication, which is bandwidth-intensive and sensitive to load imbalance.

Definition 6.11: Expert parallelism

Expert parallelism is the distribution strategy for mixture-of-experts models that places different experts on different devices and routes each token's activations to the devices hosting its selected experts through AllToAll collectives, scaling total parameter count across the cluster without scaling the compute each token consumes.

1. **Significance**: It decouples model capacity from compute budget. For illustration, a trillion-parameter MoE model may activate on the order of 10B parameters per token, so per-token FLOPs remain near those of a 10B dense model while capacity grows roughly 100×. The price is communication structure: every MoE layer requires two AllToAll shuffles per pass (dispatch and combine) whose volume scales with tokens moved times hidden dimension, stressing the cluster's bisection bandwidth rather than any single link.
2. **Distinction**: Unlike tensor parallelism, which splits dense operations so that every device computes a fraction of every token, expert parallelism is conditional computation. Each device computes only the tokens routed to its experts, which makes both the

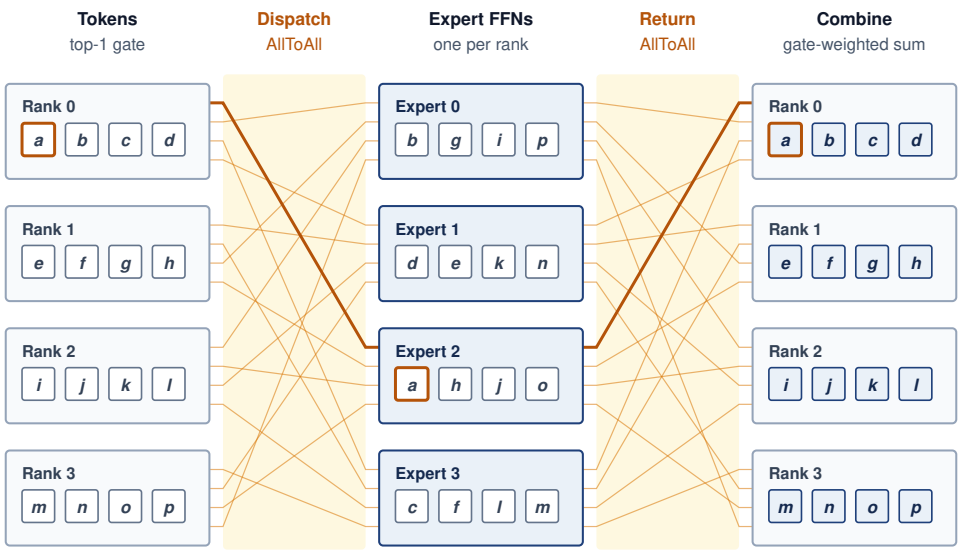

Figure 6.19: Mixture-of-Experts (MoE) AllToAll Routing: Expert parallelism on four ranks, one expert per rank, read left to right and drawn for top-1 routing with one token traced. Each rank's gate picks the top-k_{experts} experts for its tokens; the first AllToAll dispatches every token to the rank that hosts its expert, the experts compute, and a second AllToAll returns the results so each rank can form the gate-weighted sum for its own tokens. In both exchanges every rank may send to every other rank, which is why the pattern stresses the fabric's bisection bandwidth.

communication pattern and the load distribution data-dependent rather than fixed by the architecture.

3. **Common pitfall**: A frequent misconception is that experts share work evenly. A hot expert receiving, for example, 3–5× its uniform share can overflow its device's activation buffer while sibling devices idle. Capacity limits and routing penalties are possible responses; a design that retains every assignment must instead accommodate the resulting memory and scheduling demand.

↩ The terms that set an AllToAll's completion time, from traffic across the fabric bisection to hotspots created by skewed expert loads, are established in section 5.8.

AllToAll dispatch routes token activations to their assigned experts. Consider $N_{\text{experts}} = 64$ experts across 64 GPUs, processing a routing-group batch of $B = 4$ sequences of length $S = 2048$ and hidden dimension $d = 4096$. Under top-1 routing, $B \times S$ tokens each contribute one activation vector. At FP16, the logical dispatch payload is $B \cdot S \cdot d \cdot 2$ bytes, approximately 67 MB; returning expert outputs doubles the forward payload to about 134 MB per block. Local destinations reduce bytes crossing the network, multiple selected experts increase routes, and backward adds gradient transfers. The resulting service time depends on placement and load balance, including in Switch Transformer (Fedus et al. 2022) (up to 2,048 experts) and GShard (Lepikhin et al. 2021).

Uniform routing is a useful capacity baseline, not a guarantee. An expert receiving several times its uniform share, the hot expert of the definition above, needs correspondingly more token processing and activation storage. The actual skew depends on the learned router and batch. A hot expert can exhaust its buffer or delay the step while other experts remain underused, coupling memory feasibility to load balance.

One response bounds expert workload with a capacity factor C_{cap}. For top-1 routing, a capacity of roughly $C_{\text{cap}} \cdot (B \times S)/N_{\text{experts}}$ token assignments per expert provides headroom above uniform load. A dropping implementation sends excess tokens past that expert through the residual connection; implementations that retain every assignment instead need memory and scheduling to accommodate the imbalance. An auxiliary load balancing loss can encourage less concentrated routing. Models using more than one expert per token, such as Mixtral 8x7B with top-2 routing (A. Q. Jiang et al.

2024), require a capacity budget that counts all selected assignments rather than reusing the top-1 denominator unchanged.

Sparse routing gives MoE layers the same coordination problem that the embedding tables of Archetype B (DLRM at Scale) pose. In both, which device a token's data must reach depends on the data itself, so load balance cannot be fixed in advance. MoE language models nonetheless remain Archetype A (GPT-4/Llama-3) workloads, because their dense layers still carry the tensor, pipeline, and data-parallel traffic developed above.

6.5.6 Sparse embedding sharding

Expert parallelism places parameters that only some tokens touch. Recommendation models push that sparsity much further. Their parameters are dominated by embedding tables (often 10 TB+) that no single GPU can hold, and their updates are sparse, because a batch of users interacts with only a tiny fraction of the items. Sharding those tables is model parallelism over stored state rather than over computation. The parameter server (PS) architecture, formalized by Li et al. (2014), remains a common way to do it. A parameter server architecture separates workers, which compute gradients, from servers, which store parameters and apply updates.

For dense models such as ResNet or BERT, that separation concentrates traffic, as section 6.3.1.3 noted. With N workers, a server's inbound links must absorb N gradient streams at once. This incast problem is why dense training moved to ring AllReduce, in which every worker sends and receives at its own link rate and the aggregate load spreads across all nodes.

Sparse embedding tables reverse that balance. In a sparse parameter-server design, the tables are sharded across servers, and workers request only rows indexed by their local minibatches. A sparse pull returns those embedding vectors; a sparse push sends updates for the touched rows. This pattern avoids exchanging the complete table, but hot rows, lookup rates, and network contention can still limit throughput. Embedding-sharding frameworks such as TorchRec also support accelerator-based layouts, so sparse partitioning should not be equated with a CPU-only server architecture. Recommendation training at this scale is the volume's Archetype B (DLRM at Scale) regime, and it differs from the dense running example in which resource runs out first.

Whether the shards live on servers or in accelerator memory, each table must be split along some axis. Row-wise sharding partitions embedding tables by row, placing the whole embedding vector of the entity with identifier id on the shard that equation 6.3 selects:

$$s = \text{hash}(\text{id}) \bmod N_{\text{shards}} \tag{6.3}$$

Each shard s contains approximately $N_{\text{entities}}/N_{\text{shards}}$ embeddings, where N_{entities} is the total number of entities and N_{shards} is the shard count. Column-wise sharding instead partitions each embedding vector $\mathbf{e}$ across devices, as equation 6.4 writes it:

$$\mathbf{e} = [\mathbf{e}^{(0)}, \mathbf{e}^{(1)}, \ldots, \mathbf{e}^{(N_{\text{shards}}-1)}] \tag{6.4}$$

Each device stores a slice of every embedding. Hybrid sharding combines both approaches, and the choice among the three depends on lookup patterns and communication overhead (table 6.6).

Table 6.6: Embedding Sharding Strategies: Comparison of row-wise, column-wise, and hybrid embedding placement strategies, balancing lookup locality against network communication collectives and memory load balancing.

Sharding strategy	Lookup pattern	Communication	Best for
Row-wise	Single device per lookup	AllToAll gather	Uniform access patterns
Column-wise	All devices per lookup	AllGather	Hot embeddings
Hybrid	Varies by embedding	Mixed	Production RecSys

The trade-off is locality against load balance. Row-wise sharding keeps each embedding vector whole on one device, while column-wise sharding spreads its dimensions across devices. Both require returning the pieces each consuming computation needs, but placement alone does not prescribe a collective; the runtime can use remote procedure calls to the shards or collectives

according to where inputs and outputs reside. A row-wise layout commonly realizes its lookups as an AllToAll of whole vectors, and a column-wise layout as an AllGather of slices of each vector (figure 6.20). Hybrid placement can replicate or column-partition hot embeddings while row-sharding the cold tail, trading additional storage and communication against concentrated lookup load.

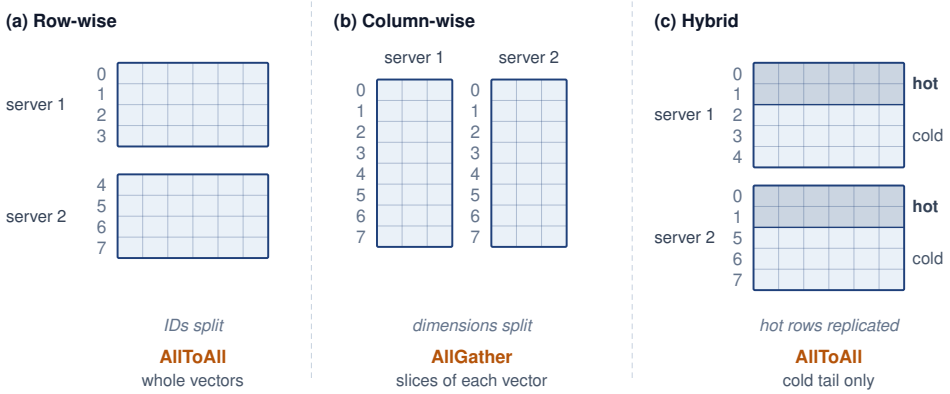

Figure 6.20: Embedding Sharding Strategies: Row-wise sharding places complete embedding vectors on specific servers based on entity ID, requiring a network gather for lookup. Column-wise sharding splits each vector across all servers, allowing parallel local lookups followed by an AllGather, which is efficient for popular "hot" embeddings. Hybrid sharding combines these approaches, using replication or column sharding for hot items and row sharding for the "cold" long tail to balance load and memory.

On accelerators, a row-wise layout turns each training step's lookups into an AllToAll exchange. Every worker sends the indices its minibatch touches to the shards that own them and receives the embedding vectors back, and the backward pass returns gradients for the touched rows along the same routes. The exchange obeys the AllToAll bound of section 5.8, and like MoE dispatch its volume and balance depend on which rows the data touches, so hot rows overload their owners unless the layout replicates or splits them.

Lighthouse 6.1: Archetype B (DLRM at Scale): DLRM vs. LLM scaling

Each workload scales through a distinct mechanism and is limited by a different resource. Dense LLMs scale through data, tensor, and pipeline parallelism, constrained by compute throughput and by the bandwidth that dense gradient and activation collectives need over fabrics such as NVLink and InfiniBand. Sparse recommendation models such as DLRM scale by sharding embedding tables of 10 TB or more across parameter servers or accelerator memory, where aggregate memory capacity, lookup rate, and the latency of many small random accesses dominate. The distinction shapes cluster design, favoring accelerator pods built around high-bandwidth dense fabrics for LLMs and memory-capacity-rich layouts for recommendation systems. This holds whether the embedding shards live in CPU-hosted servers, in accelerator HBM, or in a tiered hybrid that keeps hot rows on accelerators and cold rows in host memory.

6.6 Hybrid Parallelism

The 175B-parameter running example has now exhausted each parallelism family on its own. Data parallelism with ZeRO-3 makes the state fit, but the per-step parameter regather priced in section 6.5 would take about 3× as long as the computation it serves, while tensor parallelism confined to one 8-GPU node still leaves several hundred gigabytes of training state per device. Because replication and state sharding (section 6.3) distribute examples and storage whereas tensor and pipeline partitioning (section 6.5) distribute layer computation, **hybrid parallelism** composes both families orthogonally. Model parallelism splits the architecture so each replica fits available HBM and stays on fast links,

while data parallelism scales throughput across multiple model replicas to process 300 billion tokens within calendar limits. This orthogonal composition defines Archetype A training at frontier scale.

Lighthouse 6.2: Archetype A (GPT-4/Llama-3): Physics of 3D parallelism

Archetype A (GPT-4/Llama-3) illustrates the motivation for hybrid parallelism. Parameter count (P) must first be converted to a complete memory budget before comparison with device capacity (M_{device}), while dataset size (D) and the deadline determine the throughput target. Tensor parallelism partitions layer computations within a suitable communication domain; pipeline parallelism assigns consecutive layers to stages; and data parallelism distributes examples from D across replicas. A design combines the dimensions that address its constraints; efficient execution does not require all three in every case.

To see why the combination matters, consider four parallelism configurations for training the 175B model on a 1,024-GPU cluster organized as 128 nodes of 8 GPUs each. The progression from pure data parallelism through tensor and pipeline parallelism shows which axis makes the job fit and what each axis adds to the step time.

All four configurations run the same modeled training step, the one this book uses whenever it quotes a step time for this model on this cluster, the reference job. One step is one optimizer update over a global batch of 1,536 sequences of 2,048 tokens in 16-bit mixed precision, and each replica feeds its share of the batch through the model one sequence per microbatch. Each GPU computes at 45 percent of its peak before communication, and up to 85 percent of the data-parallel gradient reduction hides behind backward computation. Because the model, batch, and cluster are fixed, the average computation per GPU is the same 7.25 seconds per step in every configuration. In B, C, and D every replica receives the same whole number of sequences, so every GPU carries exactly that load. A is the exception, because its 1,536 sequences do not divide evenly over 1,024 replicas (1.5 each). Some GPUs would process two sequences while others process one, so the walk quotes no step time or efficiency for A. The configurations differ only in the training state each GPU must hold and in the time each adds to that computation for tensor-parallel exchanges, exposed data-parallel transfer, and the pipeline bubble.

In Configuration A, pure data parallelism (DP-1024) serves as a counterfactual rather than a feasible placement for this model. Every GPU would hold the full 2,800 GB of training state against 80 GB of HBM. The FP16 weights alone occupy 350 GB, so an A100 or H100 cannot replicate them; sharding only optimizer states does not change that fact. Communication is not what rules the layout out. If replication were possible, all 1,024 ranks would AllReduce the 350 GB gradient once per step. A hierarchical reduction first reduces inside each node over NVLink, so each GPU's InfiniBand port carries $1/8$ of the gradient, and the exchange takes about 3.1 seconds, of which about 0.46 seconds stays exposed beside the average 7.25 seconds of computation. The scaling-efficiency notebook in section 6.4 reached only 15 percent for this placement because it priced the exchange as a flat ring on each GPU's port behind a shorter FP8 compute window, which hides far less of the transfer.

Configuration B introduces intra-node tensor parallelism alongside data parallelism (TP-8, DP-128). Within each of 128 model replicas, 8 GPUs use tensor parallelism over NVLink. The corresponding tensor shard from every replica forms one 128-rank data-parallel group, so the layout contains 8 such groups. Each rank AllReduces its $1/8$ gradient shard, 43.75 GB rather than 350 GB, in about 1.74 seconds. Tensor parallelism is not free, however. Its forward and backward activation AllReduces add about 0.90 seconds per step on NVLink, so the modeled step grows to 8.42 seconds, a scaling efficiency of 86 percent. The layout still does not fit. At 16 bytes per parameter each GPU would hold about 350 GB of training state, so tensor parallelism confined to one node is not enough by itself, and the layout would also need gradient and optimizer sharding across its data-parallel group before it fits.

Configuration C adds pipeline stages across nodes (TP-8, PP-4, DP-32) so that each of 32 model replicas uses 8-way tensor parallelism within a stage and 4 pipeline stages across nodes. For each of the 8×4 tensor-and-pipeline shard positions, the corresponding ranks from all replicas form one 32-rank data-parallel group. Each rank therefore AllReduces a $350/(8 \times 4) \approx 10.94$ GB gradient shard, in about 0.42 seconds. Pipeline communication forwards activations between stages and has

lower volume than a full-model AllReduce. The trade-off is the pipeline bubble. At the beginning and end of each training step, some pipeline stages are idle while waiting for activations from earlier stages or gradients from later stages. The bubble fraction is approximately $(N_{\text{PP}}-1)/(m+N_{\text{PP}}-1)$, so with 4 pipeline stages and the 48 microbatches each replica receives per step, the bubble idles $3/51 \approx 5.9\%$ of the pipelined step, about 0.45 seconds. The pipeline cuts the data-parallel traffic by another factor of 4, but the bubble costs more time than the smaller transfer saves, so the modeled step lengthens to 8.67 seconds and efficiency falls to 84 percent.

Configuration C does not fit memory by itself either. Its training state is about 87.5 GB per GPU, above the 80 GB of an H100. Adding ZeRO-1 across its 32-rank data-parallel group closes that gap. Each GPU then keeps its tensor-and-pipeline slice of weights and gradients but only a $1/32$ share of that slice's optimizer state, about 24 GB in all. ZeRO-1 matches plain DDP's ideal data-parallel byte volume at matching communication widths (table 6.3), so the model prices the same step for it, but equal bytes do not preserve the overlap schedule. Gradient reduction to owners must precede their optimizer updates, and the updated-parameter AllGather follows those updates. Configuration C with ZeRO-1 passes this simplified persistent-state memory budget, and its predicted step holds only if that reduction, update, and parameter-gather sequence achieves the modeled exposed communication time.

Configuration D deepens the pipeline instead of sharding the optimizer (TP-8, PP-8, DP-16). It is the layout of the reference job, in which tensor parallelism spans one node and the pipeline is the shallowest power of two at which unsharded mixed-precision Adam state fits. Each GPU holds about 43.8 GB of training state, which leaves about 36 GB of the 80 GB budget for activations, workspaces, and communication buffers. Each of the 16 replicas runs its 96 sequences as 96 one-sequence microbatches, so the bubble idles $7/103 \approx 6.8\%$ of the pipelined step, about 0.53 seconds, and the 16-rank data-parallel groups reduce their smaller shards in about 0.21 seconds. With tensor-parallel, data-parallel, and bubble time all charged, the modeled step takes 8.71 seconds, a scaling efficiency of 83 percent. Against configuration C with ZeRO-1, the deeper pipeline costs about 0.04 seconds per step, its larger bubble less its smaller transfer. Its optimizer still waits for the reduced gradients, but each GPU holds the optimizer state for its whole tensor-and-pipeline slice and updates that slice in place, so the post-update parameter AllGather that ZeRO-1 adds never enters the step.

Within this step model, communication does not separate the layouts. B, C, and D finish within 3 percentage points of one another, and each added axis lengthens the modeled step slightly because it brings its own exchange or idle time. Memory decides the design. Tensor parallelism inside the NVLink domain divides the training state by the node size at the price of frequent exchanges on the fastest link, pipeline stages divide it again at the price of a bubble, and data parallelism leaves only small gradient shards for InfiniBand. Peak memory, kernel shape, and scheduling still need verification before that prediction becomes a runnable design. The infrastructure hierarchy therefore dictates the parallelism hierarchy. Here, hierarchy-aware parallelism means matching exchange requirements to measured communication domains rather than choosing each degree independently.

Definition 6.12: Hierarchy-aware parallelism

Hierarchy-aware parallelism is the strategy of mapping different parallel execution modes to the physical bandwidth tiers of the cluster.

1. **Significance**: It places high-frequency synchronization (for example, tensor parallelism) on fast links such as NVLink, while assigning lower-frequency exchanges (for example, data parallelism) to slower tiers when their transfers fit the schedule. This alignment can improve scaling efficiency (η_{scaling}) by reducing exposed communication time (T_{comm}) and latency (L_{lat}); the gain depends on readiness, contention, and available overlap.
2. **Distinction**: Unlike uniform parallelism, which treats all node-to-node links as equal, hierarchy-aware strategies respect the bandwidth cliffs between die, node, and rack boundaries.

3. **Common pitfall**: A frequent misconception is that any model can be sharded across any number of nodes. In reality, if the hierarchy mapping is wrong (for example, sharding a large tensor across a slow inter-rack link), the communication time will dwarf the compute time, making the scale-out useless.

The dependence runs both ways, because a parallelism configuration also determines which traffic the fabric must carry. In the reference job's TP-8, PP-8, DP-16 layout, tensor reductions stay within each fast local domain, pipeline traffic crosses neighboring stages, and data-parallel groups synchronize corresponding shards across replicas. Rail-optimized and hierarchical fabrics can exploit that structure, but a workload mix that adds expert AllToAll traffic changes the global bandwidth requirement. Infrastructure planning must therefore evaluate the anticipated communication groups together rather than optimizing the fabric for a single layout.

↩ Rail-optimized fabrics, which cable each GPU slot position to its own rail switch, are established in section 3.5.3.

6.6.1 The 3D training loop

A 175B-parameter model trained with 3D parallelism (figure 6.2) can be traced on a fleet configured with tensor parallelism (TP) of 8 GPUs, pipeline parallelism (PP) of 16, and data parallelism (DP) of 128. This configuration uses 16,384 GPUs (8 GPUs $\times$ 16 $\times$ 128) organized into a hierarchy of bandwidth domains. That is 16× the 1,024-GPU reference cluster, so every level of the hierarchy is exercised even though parameter count alone would not require all three dimensions. Section 6.6.2 justifies its pipeline depth.

The training step begins at the data-parallel level. Each of the 128 model replicas receives a distinct slice of the global batch. Within each replica, the model is split across 16 pipeline stages (nodes), with microbatches flowing sequentially from the embedding layer on Node 0 to the loss calculation on Node 15. At the finest granularity, within each node, the 8 GPUs fuse into a single "super-accelerator" via TP. The transformer matrix multiplications are split across these devices; column- and row-parallel pairs retain local intermediate shards and exchange partial results at the reduction points already derived. For a 175B-scale hidden dimension, the activation payload is roughly 201.3 MB. The resulting tensor-parallel traffic is about 4.2 GB per pipeline stage per microbatch in the forward pass, or 8.5 GB including the backward pass. Across the full 96-layer replica, the forward tensor-parallel traffic is about 67.6 GB. This traffic is the highest-intensity in the system, but its exposed latency can remain small relative to compute because the chips communicate over 300 GB/s–450 GB/s of local bandwidth per direction.

The backward pass exposes dependencies between parallelism dimensions. Nodes exchange activation gradients point-to-point, roughly 201.3 MB per microbatch per stage boundary in this example. A smaller payload still stalls a stage if it arrives after the stage needs it. Data-parallel synchronization creates another dependency. Although full-model gradients occupy about 350 GB, each data-parallel group exchanges only the corresponding tensor- and pipeline-parallel parameter shard. Gradient bucketing groups ready gradients as backward computation proceeds. Which exchange limits the step depends on its readiness, service time, and overlap; data parallelism is not necessarily the final bottleneck.

The architectural imperative is *bandwidth matching*, in which each exchange's startup and transfer time must fit its available compute or overlap interval. In this mapping, frequent TP exchanges use the NVLink domain (300 GB/s per direction or higher), while PP transfers cross nodes and DP reductions overlap backward computation. Accumulation can enlarge the compute interval between DP synchronizations, subject to the global-batch constraint. An infeasible mapping leaves part or all of the 16,384 GPUs fleet waiting. Comparing service time with the schedule completes the feasibility check introduced in section 6.2.1.

↩ Bucket fusion, which groups gradients that become ready together into larger reductions so communication can begin during backpropagation, is established in section 5.10.3.

6.6.2 Hybrid-parallelism worked example

Applied to physical infrastructure, bandwidth matching turns cluster design into a placement problem, and memory capacity per device (80 GB) bounds every placement. For the 3D training loop's layout on 16,384 GPUs, built from 8-GPU NVLink nodes, this reasoning leads to a concrete placement. Fix tensor parallelism at $N_{\text{TP}} = 8$ so the bandwidth-heavy matrix-multiplication collectives stay within one 8-GPU node. Map pipeline parallelism across nodes in the same high-bandwidth

rack or island, here $N_{\text{PP}} = 16$, with $N_{\text{PP}} = 8$ as the alternative when the per-device memory budget allows. Use data parallelism for the remaining scale-out dimension across pods. The result is not a generic rule but a consequence of matching each traffic pattern to the cheapest fabric tier that can carry it. With the placement fixed, the next question is whether the static model state fits inside each accelerator's HBM budget.

The memory budget for a 175B model must include every persistent training component. Under the mixed-precision Adam convention used here, each parameter has FP16 weights and gradients, an FP32 master weight, and two FP32 optimizer moments, totaling 16 bytes before activations and workspaces. With tensor parallelism degree 8 alone, these components occupy 350 GB per device, exceeding H100-class capacity. Adding 16 pipeline stages spreads that state across the product of the two degrees, yielding 21.9 GB per device under even partitioning. A budget that counts only weights and optimizer moments, about 13.7 GB per device here, understates this total. Activation peaks, layer imbalance, and communication workspaces must fit in the remaining memory before this configuration is feasible. Activation checkpointing can bring those peaks within budget, but the plan then pays in useful throughput, because recomputation costs about one extra forward pass per step when full and a few percent when selective (section 6.3.1.2), work that HFU counts and MFU does not.

Once the configuration fits in memory, its traffic still follows the same tiers, and each data-parallel group reduces only its tensor-and-pipeline shard of the 350 GB gradient buffer once per update. Gradient bucketing overlaps that reduction with backward computation, so its exposed cost depends on the computation remaining before the optimizer needs the reduced shard.

The last cost is scheduling rather than bandwidth, because a pipelined model is only efficient when enough microbatches are in flight to keep all stages busy. This remaining cost is the pipeline bubble developed in section 6.5.2. Applying its bubble fraction $\frac{N_{\text{PP}}-1}{m+N_{\text{PP}}-1}$ to this hybrid layout, $N_{\text{PP}} = 16$ stages with $m = 32$ microbatches waste $\frac{15}{47} \approx 31.9\%$ of theoretical compute capacity, nearly one-third. Configuration C's 4 stages at the same microbatch count would waste roughly 9 percent.

The deeper pipeline is chosen at this scale for two reasons. Its 16 stages spread the full training state to about 21.9 GB per device without optimizer sharding. Filling 16,384 GPUs with configuration C's 32-GPU replicas would instead need 4× as many data-parallel replicas, which multiplies the global batch $N_{\text{DP}} \times m \times B_{\text{micro}}$ by the same factor at the same per-replica schedule. The shallow pipeline spends global batch; the deep one spends bubble. The 1F1B scheduling that reclaims activation memory, and the larger microbatch counts it makes practical, drive the bubble fraction down asymptotically. Every added microbatch also enlarges the global batch, so the bubble target must be reconciled with the critical-batch constraint from section 6.4.6.

Blackwell, the accelerator generation after the H100, shows how a wider fast communication domain changes 3D parallelism choices. NVLink 5 provides 1.8 TB/s of bidirectional bandwidth per GPU, twice the H100's capacity, which shortens the transfer term for a fixed payload but does not reduce message startup latency. Larger tensor-parallel (N_{TP}) groups, such as $N_{\text{TP}} = 16$ or $N_{\text{TP}} = 32$ across nodes, remain subject to topology and local-kernel efficiency. The 10 TB/s internal link joins Blackwell's two dies behind one logical GPU, so those dies are not separate tensor-parallel ranks.

For a larger 1 trillion-parameter model on this hardware, one candidate uses $N_{\text{TP}} = 16$ across two 8-GPU nodes joined by NVLink Switch and $N_{\text{PP}} = 8$. Relative to the H100 layout's $N_{\text{TP}} = 8$ and $N_{\text{PP}} = 16$, this preserves the model-partition product while halving stage count, and the bubble must be recomputed from the new stage and microbatch counts rather than assumed to halve. Improving local bandwidth shifts pressure outward. Once intra-node exchanges limit the step less, the inter-node fabric that carries pipeline and data-parallel traffic, together with any expert AllToAll service, becomes the term to check against the next tier. The same placement constraints define a search over the TP, PP, and DP factorizations of a fleet.

Example 6.1: Automated design space search

Scenario: An engineering team schedules a 175B-parameter model across 8,192 H100 GPUs (80 GB HBM each), 8× the 1,024-GPU reference cluster so that the pipeline and data-parallel

degrees both have room to vary. The team navigates the multi-dimensional 3D-parallelism search space (TP × PP × DP).

Diagnosis: Candidate configurations must fit within 80 GB HBM and satisfy the modeled communication budget. The search evaluates TP × PP × DP factorizations and selects TP=8 GPUs, PP=8, and DP=128 as the best predicted MFU among configurations feasible under its assumptions. That result is a candidate for measurement, not proof of the best attainable runtime.

Systems lesson: This search favors TP=8 GPUs within the reference topology's fast domain and PP=8, with 36.4 percent projected MFU. The result balances memory, local computation, transfers, and pipeline idle time within the search's cost model. A different topology, microbatch schedule, or kernel efficiency can change the preferred factorization.

Adjusting any one dimension shifts the pressure onto the other two, making the factorization TP × PP × DP a tightly coupled system rather than three independent knobs. A projected MFU such as this search's still needs comparison with measured execution, and published runs report MFU well below peak that can fall as the same model moves to more GPUs (figure 9.2 in section 9.2.2).

Checkpoint 6.4: Hybrid 3D parallelism

Verify your understanding of how parallelism strategies combine:

- ☐ **Dimension roles**: identify which dimension in a TP=8, PP=16, DP=128 configuration partitions operations within each layer inside one node.
- ☐ **Bubble vs. depth**: determine how moving from PP=16 to PP=8 changes the pipeline bubble fraction.
- ☐ **DP vs. PP**: explain the trade-off, with the budget fixed at that layout's 16,384 GPUs, between increasing DP (data parallelism) and increasing PP (pipeline parallelism).
- ☐ **Latency sensitivity**: identify which dimension, TP, PP, or DP, is most sensitive to inter-node latency.

The four configurations and the worked layout tell one story. Pure data parallelism would replicate 2,800 GB of training state on every GPU, which no fabric can repair. Keeping tensor parallelism inside the NVLink domain cuts each rank's state and gradient shard by the tensor-parallel degree at the price of tensor-parallel exchanges, and pipeline stages cut them again at the price of a bubble. Communication alone does not make a layout feasible, however. Each layout must also hold the full training state, which is why configuration B would need gradient and optimizer sharding across its data-parallel group, why configuration C needs ZeRO-1, and why the reference job's TP-8, PP-8, DP-16 layout deepens the pipeline until unsharded state fits, at a modeled 83 percent scaling efficiency. The worked layout at 16× that scale goes deeper still, until the persistent state falls to about 21.9 GB per device. Hybrid parallelism is therefore the joint solution of a memory constraint, a bandwidth hierarchy, and a batch-size limit, and every axis added to satisfy one of them tightens another. Alignment training tests that joint solution hardest, because it asks one fleet to host several models with different memory and compute profiles at the same time.

6.7 Multi-Model Training for Alignment and RLHF

A cluster partitioned to train a single model cannot run post-training alignment in place, because reinforcement learning from human feedback (RLHF) coordinates up to four distinct models inside a single training step. In proximal policy optimization (PPO), an active policy generates responses token by token, a preference-trained reward model scores them, a frozen reference copy constrains policy drift, and a learned value model estimates expected returns to guide the gradient update. Those roles carry sharply different memory footprints and alternate between memory-bandwidth-bound autoregressive generation and compute-bound backward passes. Direct preference optimization (DPO) sidesteps online generation and drops the separate reward and value models, trading

live exploration for a simpler two-model loop. Designing an alignment cluster therefore starts by tallying which models, states, and execution phases must coexist on the fleet.

6.7.1 The multi-model coordination problem

Holding four models in accelerator memory at once multiplies the static state budget before a single prompt is even decoded, and only two of those models ever run a backward pass. In PPO (Schulman et al. 2017), as used in InstructGPT-style systems (Ouyang et al. 2022), the four roles differ sharply in execution mode and memory burden (table 6.7). In figure 6.21, the policy π_θ generates responses for scoring by the frozen reference and reward models and evaluation by the value model M_{value}. The penalized reward ($r - \beta_{\text{KL}}$ KL) subtracts a Kullback–Leibler divergence penalty, weighted by a coefficient β_{KL} unrelated to link bandwidth β, to discourage drift from the reference. Generalized advantage estimation (GAE) combines rewards and value estimates to judge how much better a response was than expected. Those estimates guide policy updates; the value model learns to predict returns, the cumulative discounted reward that follows each generated token. Reference and reward models remain forward-only in this loop, and shared-backbone implementations can combine roles without changing their distinct responsibilities.

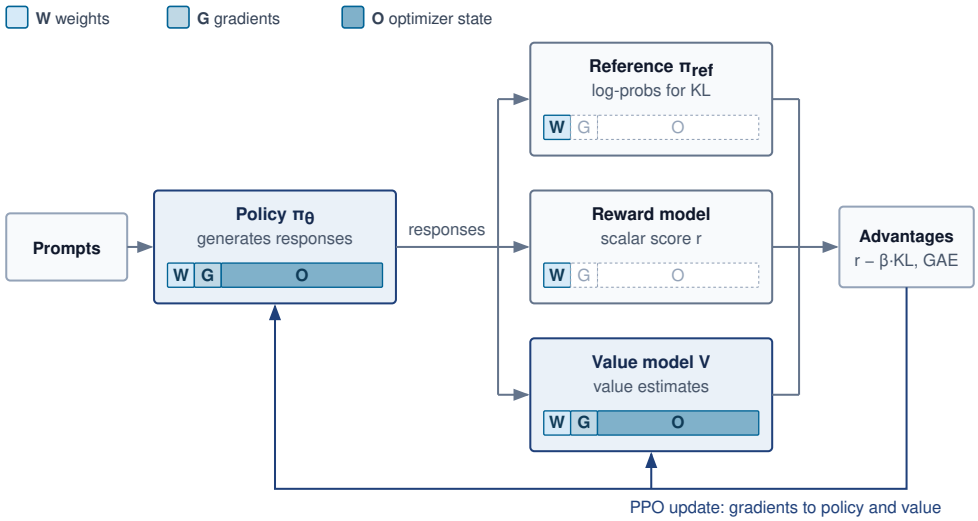

Figure 6.21: One PPO-Style RLHF Step: The policy generates responses to a batch of prompts. The frozen reference model supplies log-probabilities for the KL penalty, the frozen reward model supplies a scalar score, and the trained value model supplies value estimates. The penalized reward and value estimates form advantages, which update the policy and value models. Blue cards carry gradients and optimizer state; slate cards run forward passes only and hold weights alone, as each card's weight, gradient, and optimizer-state strip shows; strip widths follow the 2, 2, and 12 bytes per parameter of mixed-precision Adam, unlike the equal-width columns of figure 6.7. That asymmetry drives placement.

Table 6.7: RLHF Model Coordination Roles: PPO-style RLHF coordinates models with different execution modes and memory footprints, so placement decisions must separate trainable policy and value state from inference-only reference and reward state.

Model	Role in PPO-style RLHF	Execution mode	Memory burden
Policy Model	The model being aligned and optimized.	Training mode with gradients, optimizer moments, and activations.	For a 70B-parameter model in mixed precision with Adam, memory is approximately 70B $\times$ (2 + 2 + 12) = 1,120 GB, the same budget as standard pretraining.

Model	Role in PPO-style RLHF	Execution mode	Memory burden
Reference Model	A frozen copy of the pretrained policy that supplies the KL-divergence penalty.	Inference mode only.	FP16/BF16 parameters without gradients or optimizer state; a 70B model needs about 70B $\times$ 2 = 140 GB, roughly 8× less than the training configuration.
Reward Model	A preference-trained scorer that produces scalar rewards for generated sequences.	Inference mode only.	Often smaller than the policy; a 13B reward model requires approximately 26 GB in FP16, but it must process every generated sequence.
Value Model	The PPO critic that estimates expected future reward.	Training mode, either separate or sharing a policy backbone.	A full-size separate critic can add another 1,120 GB; smaller models or shared-backbone designs reduce this to 200–400 GB.

The aggregate memory demand of the four-model PPO system dwarfs standard pretraining. A naive co-location of a 70B policy, 70B reference, 13B reward, and 13B value model requires approximately 1,494 GB of accelerator memory, before accounting for the KV caches and intermediate activations generated during sequence generation. A key-value (KV) cache stores attention keys and values for previously processed tokens so the policy can generate the next token without recomputing that state. It grows with the number and length of active sequences, creating a temporary memory budget distinct from weights and optimizer state. On H100 GPUs with 80 GB of HBM each, this system requires a minimum of 19 GPUs as an aggregate static-state capacity lower bound, including optimizer state and gradients. This division does not establish a feasible placement, because each phase must also fit its KV or activation peak, workspaces, and uneven shards. Replicated policy groups increase the total state further.

6.7.2 Infrastructure asymmetry between training and inference models

The defining infrastructure challenge of RLHF is not the *total memory footprint* but the *asymmetry* between the models' compute profiles. The policy and value models require full backward passes with gradient computation, activation checkpointing, and optimizer updates, compute-intensive operations that benefit from tensor parallelism and high arithmetic intensity. The reference and reward models, by contrast, perform only forward passes over generated sequences. Their scoring passes do not require optimizer state, and they need not decode token by token. The policy's rollout phase creates the autoregressive generation workload and its KV-cache budget. Separating scoring from generation matters because forward-only execution does not imply a single memory or compute profile.

The asymmetry creates a placement dilemma. Co-location permits devices to serve different phases, but inactive model state can occupy memory needed for the active phase. Dedicated pools allow each role to use a suitable execution configuration, but pools may idle while waiting for another role's output. Separation also adds transfers of generated sequences, scores, and updated policy weights. The design must compare phase utilization and data movement rather than assume that either placement eliminates waste.

The PPO rollout phase introduces autoregressive dependencies absent from the teacher-forced pretraining baseline, in which each input token comes from the training data rather than from the model's previous output. The policy generates responses before sequence-level rewards guide its updates. At small active batches, decoding is often memory bandwidth bound because each new token requires another model pass. Each token in the KV cache holds a key vector and a value vector per KV head in every layer. For a Llama-2-70B-style grouped-query attention policy generating 1,024 tokens for each of 256 active sequences, those per-token bytes, written out in equation 10.5, put the generated-token portion at 85.9 GB. That portion alone already exceeds the 80 GB of HBM on one H100, and prompt tokens add KV state as well. This temporary cache must fit alongside the state live during rollout; its lifetime differs from training activations and optimizer state.

RLHF systems address this asymmetry through temporal multiplexing. During the generation phase, the cluster behaves like an inference system: the policy model generates sequences while the reference model computes log-probabilities, and batching active sequences alongside their KV state determines how much useful work shares each weight read. During the training phase, the same

↪ The capacity wall this cache imposes on serving systems is developed in section 10.7.1.

fleet switches back to training mode and computes gradients through the policy and value models using the 3D parallelism configuration of section 6.6, where activation residency and gradient synchronization again govern the schedule. Because PPO alternates between inference-shaped rollout work across four model roles and training-shaped updates, reformulating the alignment objective itself can eliminate part of that infrastructure overhead.

6.7.3 Resource requirements for PPO versus DPO

Standard offline DPO (Rafailov et al. 2023) takes that route by removing the separate reward and value models, along with the online rollout phase, from the update loop, reformulating alignment as a supervised loss over offline preference pairs. The policy learns directly from preferred and dispreferred responses using their log-probabilities relative to a frozen reference; because those responses are already in the training data, the update requires neither autoregressive generation nor advantage estimation.

A DPO implementation that retains a resident reference has two model roles, the trainable policy and the frozen reference. For fixed preference pairs and a fixed reference, reference log-probabilities can instead be precomputed, removing that model from the repeated update loop. Offline DPO needs no autoregressive rollout during those updates. Variable lengths still require padding, packing, or length-aware batching, and policy training can use the same partitioning, accumulation, and overlap mechanisms as pretraining.

The infrastructure difference follows from how each method obtains feedback. Offline DPO learns from a fixed preference dataset, so the policy never receives feedback on new responses. PPO-style training scores responses from the evolving policy and can expose behavior the dataset lacks, provided the reward signal is reliable. The generation phase, the reward and value models, and the memory they occupy are the price of that online loop, so the choice is settled by achieved quality per unit of total resource cost.

To make the infrastructure trade-off concrete, consider aligning a 70B policy model using a policy/reference pool of 256 H100 GPUs (80 GB HBM each). PPO's separate reward and value pools sit outside that allocation. The policy is a 70B-class model rather than the 175B reference, a size at which each role's memory budget stays readable while the four PPO roles together still need at least 19 GPUs, more than one node.

Napkin Math 6.9: RLHF infrastructure budget: PPO vs. DPO

Problem: This alignment setup co-locates a 70B policy model with a 70B reference model and, for PPO, uses pools outside the 256-GPU policy/reference allocation for a 13B reward model and a 13B value model. Under 3D parallelism (TP=8 GPUs, PP=4, DP=8), how much memory resides on each policy GPU, and which separate model pools and generation-phase memory requirements does DPO remove?

PPO memory budget (per-GPU, with TP=8 GPUs, PP=4, DP=8 for policy)

The policy model under 3D parallelism with TP=8 GPUs and PP=4 distributes one replica's 1,120 GB training state across 32 GPUs, yielding 35 GB per GPU. The reference model, requiring only 140 GB for inference, is co-located with the policy GPUs at an additional 4.4 GB per GPU (140 GB / 32 GPUs). It could instead be sharded across an additional 8 GPUs pool outside the policy allocation at 17.5 GB each. The 13B reward model adds 26 GB shared across its pool. The 13B value model in training mode adds approximately 208 GB (13B $\times$ 16 bytes/param) across its pool.

Total static memory (one co-located policy + reference replica on 32 GPUs): 35 GB + 4.4 GB $\approx$ 39.4 GB per GPU, leaving $\sim$ 40.6 GB for activations and KV cache. With generation-phase KV caches for a rollout batch containing 256 sequences and 1,024 tokens inside one $TP \times PP$ policy replica, each policy GPU must reserve an additional 85.9 GB/32 GPUs $\approx$ 2.7 GB for its KV cache shard. The remaining $\sim$ 37.9 GB is generation-phase headroom. Once rollout KV state is released, the update phase recovers the 40.6 GB headroom before allocating training activations and other runtime buffers. The phase peaks must be evaluated separately.

DPO memory budget (same parallelism configuration)

The policy model uses the same 35 GB per GPU. The reference model adds 4.4 GB per GPU (co-located). No reward model, no value model, no KV cache for generation.

On the policy GPUs, total static memory is 35 GB + 4.4 GB = 39.4 GB per GPU, identical to PPO's co-located policy/reference footprint. Both have 40.6 GB available before update-phase activations and runtime buffers when PPO releases its rollout KV cache. DPO removes the separate generation peak and the reward/value model pools; the absence of rollout KV state alone does not establish larger update microbatches.

Throughput comparison

PPO alternates generation and parameter updates. If generation consumes 60 percent of an iteration in this illustrative schedule, updates occupy the remaining 40 percent. That time fraction is a share of the schedule rather than an MFU, so update-phase MFU must be multiplied by it to report update computation over the whole iteration, using the same allocated hardware denominator. DPO eliminates online generation, but its measured throughput still depends on sequence lengths, reference scoring, parallelism, and memory use. The static-state comparison in table 6.8 is therefore kept separate from these schedule-dependent effects.

Table 6.8: PPO vs. DPO State and Execution Phases: Static-state totals sum one copy of each model across its device pool, including trainable optimizer state and gradients. Data-parallel replication and phase-specific runtime allocations add to these totals; the table does not determine a minimum device count or measured MFU.

Metric	PPO (70B policy)	DPO (70B policy)
Models required	4 (policy, ref, reward, value)	2 (policy, reference)
Static state, one copy of each model	$\sim$ 1,494 GB	$\sim$ 1,260 GB
Device count	Depends on placement and phase peak	Depends on placement and activation peak
Generation phase	Yes (autoregressive)	None
Update throughput	Shares the schedule with generation	No online-generation phase
Data pipeline	Online generation	Static preference pairs

Systems insight: In this 13B reward/value-model scenario, DPO reduces the one-copy static-state budget by 15.7 percent by eliminating the reward and value models. It also removes online generation, but the throughput benefit requires a measured schedule. Infrastructure constraints therefore weigh on the choice between PPO and DPO as much as modeling preferences do. An organization with a limited GPU budget may prefer DPO because it restores the workload to a fixed-data training loop. If the PPO value model is policy-sized rather than 13B, the reduction is larger still, from 2,406 GB to 1,260 GB.

6.7.4 Sequence length variance and batching challenges

The throughput comparison in notebook 6.9 noted that even after DPO removes online generation, realized step time still depends on sequence-length variance. Pretraining often packs documents into fixed-length blocks, such as 2,048 or 4,096 tokens, whereas alignment batches combine prompts and responses whose lengths can span a 10–50$\times$ range within a single minibatch. Those variable lengths directly govern activation residency, padding waste, and load balance across workers.

Length variance creates two interacting problems. Padding every sequence to the longest in a batch can leave most token positions unused. A batch of one 2,048-token sequence and fifteen 128-token sequences, padded to 2,048 tokens, fills only 3,968 of its 32,768 positions. Kernels that execute padded positions therefore spend about 88 percent of their work on padding, while kernels that skip them recover most of that waste. Dynamic batching groups similar lengths to reduce padding, but assignments across workers still need balancing. A worker processing longer sequences can have a larger activation peak and finish later, leaving other workers waiting for the required gradient contributions.

Sequence packing combines short sequences within a fixed-length input, with attention masking that keeps their contexts separate, and adaptive microbatching sizes each microbatch by token count or predicted memory. Packing must also respect the weighting rule of section 6.3.1.1. The reduction must keep the objective's own denominator, valid-token counts for a token-mean PPO loss and pair counts for standard DPO (Rafailov et al. 2023), so packing changes layout but not weighting. Balancing token work across workers then remains a performance goal rather than a correctness requirement.

RLHF shows why no single parallelism strategy is sufficient by itself. The policy model may need tensor and pipeline parallelism, the reference and reward models behave like inference services, the value model may require training state, and the rollout data pipeline changes shape from step to step. The procedure that handles this case is the general one, which identifies the limiting constraint first and then chooses the parallelism pattern that moves the least data while fitting the required state.

6.8 Parallelism Strategy Comparison

A parallelism strategy is useful only if it moves the limiting constraint without creating a larger one. Assembling data, tensor, pipeline, and hybrid partitioning into a coherent decision procedure requires evaluating where each strategy relocates overhead across the cluster's memory and interconnect hierarchy. For a new model, the decisive questions are which memory state fits, how often communication occurs, where the pipeline bubble appears, and which abstraction cost the engineering team can tolerate. Sequence, context, and expert parallelism sit beside the 3D core in table 6.9, and they remain subject to the same batch, memory, and topology constraints.

Table 6.9: Parallel Training Strategies: Data, tensor, and pipeline parallelism form the canonical 3D core. Sequence/context parallelism extends the partition along token positions, while expert parallelism extends it across sparse model capacity. Strategy selection remains an engineering feasibility test: identify the binding constraint, the communication primitive it creates, the placement decision it forces, and the measurement that proves the strategy is helping.

Strategy	Binding constraint	Primary movement	Placement decision	Validation signal
DDP	Model state fits on one accelerator	AllReduce gradients every step	Replicate full model across data-parallel workers	Step time dominated by compute, not synchronization
FSDP/ZeRO	Model state exceeds one accelerator	Gather parameters; ReduceScatter gradients	Size prefetch and reshard windows within peak memory	State plus activations and workspaces fit
Tensor parallelism	One layer exceeds local capacity or runs too slowly	Reduce or gather partial layer results	Keep tightly coupled partitions on fast local links	Exposed collective time fits the layer-time budget
Pipeline parallelism	Layer stack fits only when split by stage	Activations and gradients between stages	Balance stages and live microbatch activations	Bubble and peak activation memory meet budget
Sequence/context parallelism	Activations or attention context exceed local capacity	Exchange sequence shards or attention state	Partition long sequences across a low-latency group	Longer context fits without communication dominating
Expert parallelism	Sparse expert capacity exceeds local memory	Route token activations with AllToAll	Balance expert load across a high-bisection group	Expert imbalance and dropped-token rate stay bounded
Hybrid parallelism	More than one constraint binds at once	Multiple collectives on coupled axes	Satisfy divisibility, batch, topology, and failure domains	MFU improves within memory and network budgets

A decision path condenses the table, starting from memory fit and then asking which structure breaks the budget (figure 6.22). The 175B reference model takes the layer branch on step time rather than capacity. An average layer holds about 29.2 GB of training state, which fits one GPU. Without tensor parallelism, however, the rule that sets configuration D's depth, the shallowest power-of-two pipeline at which unsharded state fits, gives 64 single-GPU stages instead of 8. At the same 96 microbatches per replica, that pipeline's bubble idles 39.6 percent of each step rather than 6.8

percent, and the modeled step grows from 8.71 seconds to 12.04 seconds. Tensor parallelism inside the NVLink node is what keeps the pipeline shallow, and because its 350 GB shards still exceed one GPU, the path ends at combining axes. Before accepting any branch, check peak memory, global batch, communication service time, and load balance together; bandwidth is a feasibility constraint throughout the decision, not an afterthought.

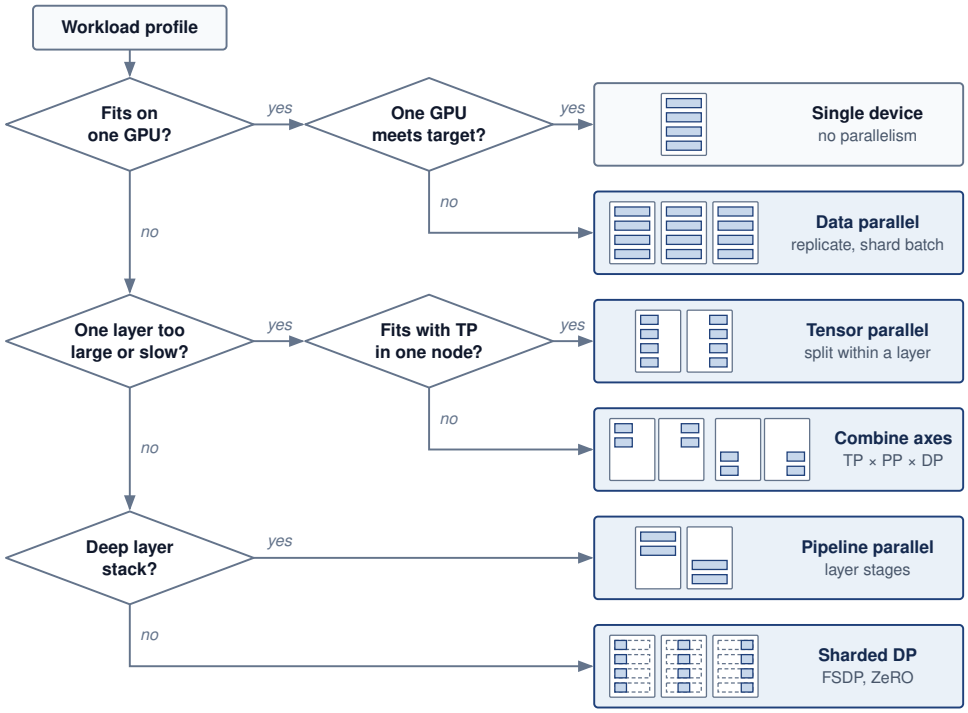

Figure 6.22: Parallelism Strategy Decision Framework: A selection guide that begins with memory fit. For a model that fits on one device, evaluate data parallelism against the throughput target and input-delivery budget; the complete dataset need not fit in accelerator memory. A model that does not fit first asks whether a single layer is too large or too slow for one device. If so, tensor parallelism splits each layer inside one node, and a model whose tensor-parallel shards still exceed device memory, such as the 175B reference model, combines tensor, pipeline, and data parallelism. A model whose layers fit is routed by depth, with a deep stack leading to pipeline parallelism and otherwise sharded data parallelism (FSDP/ZeRO). Any outcome can add data parallelism for throughput, and peak memory, global batch, bandwidth, and workload imbalance constrain every candidate.

Identifying the limiting constraint comes before choosing an implementation stack, which is why both the strategy comparison and the decision path stay deliberately coarse. Dense large-model training and recommendation-style workloads motivated each of these partitioning strategies. The same constraint logic also reaches the federated edge case (table 6.10), where the bottleneck shifts from accelerator memory and fabric bandwidth to unreliable devices, privacy, and intermittent connectivity. A configuration chosen across these regimes still has to be expressed in software, and the way a framework expresses it decides which of the chapter's budgets the engineer controls directly.

Lighthouse 6.3: Distributed archetype spectrum

Partitioning and coordination choices differ with the system's primary bottleneck:

Table 6.10: Partitioning Strategy by Archetype: Partitioning and coordination choices differ with the dominant bottleneck.

Archetype	Primary partitioning strategy	The logic
Archetype A (GPT-4/Llama-3)	Hybrid 3D parallelism	Combine tensor (width), pipeline (depth), and data (throughput) parallelism to fit roughly 3.5 TB of weights under third-party estimates (OpenAI has not disclosed the parameter count).
Archetype B (DLRM at Scale)	Embedding sharding	Partition massive 10 TB+ tables across parameter servers or accelerator shards; use sparse AllToAll updates.
Archetype C (Federated MobileNet)	Federated learning	The same logic extends to coordinating on-device MobileNet updates across unreliable, privacy-constrained edge devices; section 11.7 develops this regime.

6.9 Frameworks and Runtime Contracts

A TP-8, PP-16, DP-128 layout that passes every feasibility check on paper can still hang at its first collective if two ranks disagree on the membership of their communication group, or run slowly if gradient buckets become ready in an order the collective runtime cannot overlap. Choosing DDP, FSDP/ZeRO-style sharding, tensor parallelism, or pipeline parallelism is therefore also a decision about which constraints the runtime owns and which constraints the engineer must manage directly. Sparse workloads add another contract, the parameter-server push/pull interface (Li et al. 2014) developed for Archetype B in section 6.5.6.

The training framework defines process groups (its records of which ranks form each tensor-parallel, pipeline, or data-parallel group), gradient buckets, and the point in backpropagation when each reduction becomes ready; the collective runtime maps those reductions onto ring, tree, or hierarchical algorithms. The engineer must choose compatible group membership and bucket order and verify which communication remains exposed on the step's critical path.

Production frameworks expose these same strategies under different names, which makes a translation aid useful when reading their documentation or configuration files. Four widely used stacks name the chapter's strategies differently (table 6.11); the names change between releases, but the partition each one implements does not.

Table 6.11: Framework Names for the Parallelism Strategies: A translation aid, not a ranking. Each entry names the mechanism a framework uses for the partition developed in this chapter; exact class and option names vary across releases.

Strategy	PyTorch	DeepSpeed	Megatron-LM	JAX
Data parallel	DistributedDataParallel (DDP)	DeepSpeed engine, ZeRO stage 0	Data-parallel process groups	Batch axis sharded over a device mesh
Sharded data parallel	Fully Sharded Data Parallel (FSDP)	ZeRO stages 1–3	Distributed optimizer	Parameter and optimizer state sharded over a mesh axis
Tensor parallel	DTensor tensor-parallel APIs	AutoTP or Megatron integration	Column- and row-parallel linear layers	Sharding annotations on model axes, partitioned by the compiler
Pipeline parallel	`torch.distributed.pipelining`	Pipeline engine	Pipeline schedules (1F1B, interleaved)	User-written stage schedules

The training design therefore determines message size, readiness order, and frequency, while collective algorithms determine how those messages move. No framework abstraction changes those physical costs, so distributed-training efficiency ultimately depends on interconnect bandwidth, memory capacity, synchronization latency, and the useful overlap admitted by the computation graph. Confusing the simplicity of a framework API with exemption from those physical costs produces the recurring engineering mistakes below.

↩ Overlap of gradient exchange with backpropagation, the lever that takes communication off a step's critical path, is established in section 5.10.

6.10 Fallacies and Pitfalls

Even when a framework makes switching from single-node DDP to multi-node execution a one-line configuration change, the underlying physics still governs the step time. In the GPT-2 scale-out of notebook 6.1 and notebook 6.2, moving the same model from one NVLink node to four Ethernet-connected nodes quadrupled the GPU count but cut parallel efficiency from 99.4 percent to 27.9 percent. The following misconceptions generalize that error, each treating one term of the step-time accounting as if it were the whole budget.

Fallacy: *Linear speedup is achievable with sufficient engineering effort.*

Amdahl's Law limits speedup when a fixed serial component remains exposed. Synchronous gradient aggregation creates an update dependency, but some of its work can overlap backward computation. As section 6.4 shows, $\eta_{\text{scaling}} = 1/(1 + N(T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}})/T_{\text{compute}})$ depends on the exposed overhead relative to useful compute as N changes. A 1,000-GPU allocation therefore does not by itself justify a 1,000$\times$ speedup forecast. Local kernel efficiency, memory feasibility, communication, and convergence must be evaluated against the specified baseline. Linear scaling is a result to demonstrate for an operating range, not a return guaranteed by device count.

Pitfall: *Optimizing MFU without considering scaling efficiency.*

A single-node MFU measurement does not predict fleet throughput by itself. For a fixed workload and hardware peak convention, scaling efficiency compares measured distributed throughput with ideal replication of that single-node baseline. Multiplying the baseline MFU by this scaling efficiency gives distributed MFU; multiplying an already measured distributed MFU by the same scaling efficiency counts communication and synchronization losses twice. The 175B reference job makes the difference concrete. A single node running at 45 percent MFU, scaled with the job's 83 percent modeled scaling efficiency on the 1,024-GPU cluster (TP-8, PP-8, DP-16, configuration D of section 6.6, with tensor-parallel and data-parallel communication and the pipeline bubble all charged), delivers about 37 percent MFU across the fleet. A planner who measures that 37 percent on the full cluster and then applies the 83 percent again would budget for about 31 percent and buy capacity the job does not need. Capacity planning must state the device count and measurement window for each metric. Checkpointing, recovery, and replay require additional accounting only when those losses fall outside that window. Useful goodput must count retained model progress rather than repeated work, under the accounting boundaries of section 9.2.3.

Fallacy: *Hyperparameters tuned on small clusters transfer directly to large-scale training.*

Engineers tune hyperparameters on 8-GPU workstations then deploy to 256-GPU clusters expecting identical behavior. At scale, convergence behavior changes. The most sensitive hyperparameter is learning rate. As section 6.3 explains, global batch increases with GPU count when local batch and accumulation depth remain fixed, so the learning-rate schedule must be revalidated. The linear scaling rule of section 6.4.5, which the one-hour ImageNet run established, suggests $\eta_{\text{new}} = \eta_{\text{base}} \times (B_{\text{new}}/B_{\text{base}})$, but this relationship holds only within bounds. As models scale, they eventually encounter the critical batch size of section 6.4.6, where adding more data per step yields diminishing returns in convergence.

Beyond the critical batch size, this relationship breaks down in a model-, data-, and optimizer-dependent way. A team that takes a small-cluster learning rate, multiplies it mechanically with GPU count, and jumps directly to a much larger global batch may see lower final quality or slower sample efficiency even though hardware throughput improved. Warmup schedules, weight decay adjustment, layer-wise optimizers such as LARS/LAMB, and careful momentum tuning can recover accuracy in some regimes, but require systematic experimentation at target scale. Organizations that skip these scaling studies waste thousands of GPU-hours on suboptimal runs.

Pitfall: *Adding GPUs to data-parallel jobs without modeling communication.*

More GPUs improve time to target only if the gain in processing rate outweighs communication overhead and any additional sample use. With fixed local batch size, adding data-parallel workers increases the global batch; fixing the global batch instead reduces local work and can expose more communication. The critical-batch analysis in section 6.4.6 supplies a workload-dependent statistical guide, not universal GPU-count or model-size thresholds. Compare complete learning curves and cost at the same target quality before committing to a larger allocation.

Fallacy: *Memory capacity alone determines the parallelism strategy.*

A model that exceeds device capacity needs a feasible partition, but memory fit does not select the fastest one. Tensor parallelism exchanges partial results within layers, while pipeline parallelism transfers stage-boundary activations and introduces scheduling bubbles. Their relative cost depends on layer shapes, microbatch count, placement, and overlap. Compare complete training-state and activation budgets together with measured time to target; an unmeasured throughput estimate or a weights-only fit cannot settle the choice.

Pitfall: *Applying FSDP or ZeRO without measuring efficiency trade-offs.*

Engineers adopt FSDP universally after reading that it "reduces memory and enables larger models." Sharding trades memory capacity for communication overhead. Full-shard FSDP gathers parameters before layer execution, gathers them again for the backward pass if the forward pass resharded them, and reduces gradient shards during backward. The relevant comparison begins with the complete training-state budget. Under the mixed-precision Adam convention established here, the 7B worked example exceeds single-device capacity before activations, so standard DDP cannot fit that state on an A100-80 GB. A weights-only fit does not establish a feasible training baseline. When the complete state, activations, and workspaces fit without sharding, DDP can avoid the repeated gathers; when they exceed the 80 GB capacity, the comparison must include a feasible alternative such as sharding or offload. FSDP provides value when it enables training at all or when the larger local batch it permits offsets its communication cost. ZeRO-Offload extends capacity through CPU memory but adds transfers of its own. The decision therefore requires measuring peak memory and exposed communication for the same workload and optimizer contract.

Fallacy: *Parallelism overhead is roughly constant regardless of model size.*

Engineers benchmark parallelism strategies on convenient small models then apply those conclusions to large-scale training. For dense transformers at fixed tokens per update, computation and gradient bytes both grow roughly proportionally with parameter count, as the earlier scaling analysis established. Larger parameter count alone therefore does not imply a better computation-to-communication ratio. Architecture, sequence length, local batch size, partition degree, and kernel efficiency determine how much useful work each rank performs; gradient representation and process-group membership determine how much it communicates. A larger model can also require a different partition that introduces new layer collectives or pipeline bubbles. Reliable strategy selection holds the workload and target quality explicit, then recomputes memory fit, local compute time, communication volume, and the available overlap window at the target scale. A small-model benchmark becomes transferable evidence only when those governing terms remain comparable.

Pitfall: *Treating scaling efficiency as a property of the hardware alone.*

A vendor benchmark might report 85 percent scaling efficiency at 1,024 GPUs, and procurement teams then plan capacity as if that number were a cluster constant. Scaling efficiency is a joint property of four things: the workload's communication-to-computation ratio, the parallelism strategy (data, pipeline, tensor, hybrid), the network topology (fat-tree vs. torus vs. dragonfly), and the software stack (NCCL version, kernel fusion, overlap quality). Reference numbers from dense large-language-model benchmarks transfer poorly to other workloads and fabrics. They may overstate scaling for graph neural networks with irregular communication and for embedding-heavy recommendation models, whose AllToAll embedding exchanges and many small random lookups load the fabric differently from dense gradient AllReduce (section 6.5.6). They also say nothing about a commodity-Ethernet placement like the GPT-2 scale-out, which lost most of its parallel efficiency on leaving the

node. Capacity planning should use measured scaling on the target workload, not a vendor hero benchmark.

Fallacy: *Gradient accumulation is free.*

Gradient accumulation can reduce synchronization from every microstep to once per window of k microsteps, eliminating $(k-1)/k$ of those synchronization events when the implementation defers reduction. The saving has costs. A resident gradient buffer remains live throughout the accumulation window; for the 7B model, the FP16 buffer occupies 14 GB. Retaining graphs across microsteps or increasing the local batch size can add activation pressure. Accumulating eight microsteps makes updates $8\times$ less frequent per processed microstep, but wall-clock spacing also depends on the communication saved. Larger effective batches can change convergence, and narrow accumulation formats can lose small contributions or overflow. The dtype, loss normalization, and reduction schedule must therefore be validated at the intended batch size. Accumulation trades communication frequency against update spacing, memory residency, and numerical stability.

6.11 Summary

Scaling laws keep pushing frontier training budgets past what one machine can hold or finish in time, and distributed training is how the fleet absorbs that growth. It confronts the scaling wall, the workload-dependent operating point beyond which communication and coordination can consume the compute time saved by adding accelerators. A feasible design must preserve the intended optimization update while satisfying model-state, activation, batch-size, and interconnect constraints. Hardware throughput alone is insufficient; the comparison must also account for time to the target model quality.

Each partition relocates cost along the C^3 axes. Data parallelism divides compute across replicas and pays in communication, the gradient aggregation whose exposed share set the scaling efficiency of the 175B communication example. Tensor parallelism divides work inside a layer and pays in communication on the critical path, which confines it to a fast domain. Pipeline parallelism divides layer depth and pays mostly in coordination, the fill and drain bubble that only more microbatches, better balance, or a smarter schedule can shrink. The 3D parallelism cube in figure 6.2 unites these trade-offs in one coordinate space ($N_{\text{DP}} \times N_{\text{PP}} \times N_{\text{TP}}$). A feasible point must satisfy peak memory and global-batch constraints as well as communication and coordination budgets.

Sharding and model partitioning make different parts of the training state fit, while synchronization and precision policies determine whether the resulting updates remain useful. Hybrid layouts compose these choices within the bandwidth hierarchy. The alignment workload extends the same reasoning to multiple models and alternating generation and update phases, where peak memory depends on which state is live together.

The chapter's strategies are complementary dimensions of one design space rather than alternatives. Systems such as Megatron-LM map tensor, pipeline, data, and expert exchanges to the communication domains that can carry them, and that matching is what delays the communication-bound regime as the fleet grows. The workload archetypes weight those dimensions differently. Archetype A (GPT-4/Llama-3) combines dense-model partitioning with state sharding and, at long contexts, sequence or context partitioning. Archetype B (DLRM at Scale) shards sparse embedding state and pays in lookup traffic rather than layer collectives. Archetype C (Federated MobileNet) moves the coordination problem onto unreliable edge devices. In each case the limiting constraint, not the archetype label, selects the configuration.

Key Takeaways: Parallelism relocates overhead

- **Splitting work moves the bottleneck**: Distributed training does not make the underlying computation smaller; it converts memory pressure into communication volume, synchronization delay, or idle pipeline time (principle 12). The winning strategy is the split that sends overhead to the least constrained part of the fleet.
- **The slowest worker sets synchronous step time**: Under bulk-synchronous updates every step waits for its last required contribution, so a rare per-worker delay becomes a

near-certain per-step delay at thousands of workers. That wait preserves the single-device global-batch update, which also needs matching sampling, valid-item denominators, and optimizer state; stale and asynchronous protocols trade that wait for staleness, which can cost extra steps to the same target quality.

- **Throughput must serve convergence**: Larger global batches can improve hardware efficiency while yielding diminishing progress per sample. Compare time to the same target quality, and treat the critical-batch estimate as a workload-dependent guide for selecting batch size and update schedules.
- **Sharding buys memory with messages**: ZeRO and FSDP partition optimizer state, gradients, and parameters (principle 2), reducing replicated storage. Peak memory must still include activations, temporary gathers, and workspaces; exposed ReduceScatter and AllGather costs determine whether the capacity gain also improves throughput.
- **Tensor parallelism needs a fast domain**: Splitting operations inside layers creates frequent communication on the dependency path. Place those groups on links with sufficient bandwidth and low latency; the relevant boundary is the communication domain, which may span nodes, rather than the server enclosure.
- **Pipeline parallelism trades bytes for bubbles**: Layer staging reduces communication frequency, but fill and drain slots leave accelerators idle (principle 10). Microbatching with $m \gg N_{\text{PP}}$ is the mechanism that turns model depth into throughput rather than pipeline slack.
- **Hybrid layouts are joint decisions**: Combine only the partitioning and sharding strategies needed to satisfy the workload's constraints (principle 14). Model divisibility, batch size, peak memory, communication groups, and failure domains must form one feasible configuration; improving one axis can expose a limit on another.

The displacement of overhead (principle 14) makes a feasible partition the beginning of the design problem. A memory-saving split is useful only if its communication and schedule still meet the run's time-to-quality budget. Every mechanism in this chapter shrinks the per-device compute term T_{compute}/N or the per-device training state, and pays for it in $T_{\text{comm}}(N)$, in $T_{\text{sync}}(N)$ through barriers, stragglers, and pipeline bubbles, or in lost overlap. In C^3 terms, no partition removes the overhead of scale; it only relocates that overhead among compute, communication, and coordination. The engineering judgment is therefore to compare complete configurations, including their exposed dependencies, rather than optimize each parallelism axis independently.

What's Next: From logic to resilience

The fleet now has its traffic patterns (communication) and its logic (parallelism). Every configuration in this chapter, however, assumed that every rank delivers its contribution on time and with correct data. A collective is not a barrier, but no dependent consumer can proceed until every contribution arrives, so one crashed node, failed link, or GPU that slows down without failing stalls the whole synchronous group. A contribution corrupted in transit is reduced into every replica's update. Across thousands of accelerators such events stop being rare. Chapter 7 begins by computing how often they arrive, then builds the checkpointing, recovery, and elastic training that keep the computation making progress on imperfect hardware.

7

Fault Tolerance

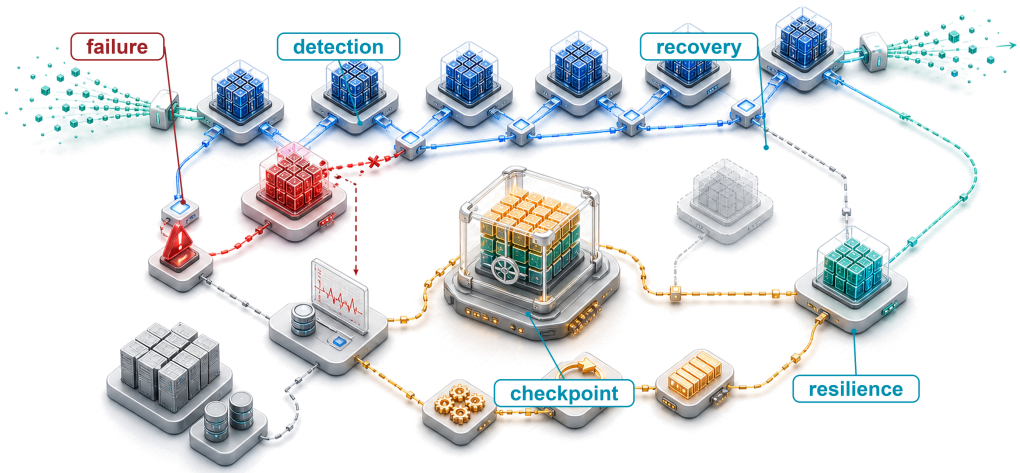

7.1 Failure Analysis at Scale
7.2 The Mathematics of Inevitable Failure
7.3 Failure Taxonomy
7.4 Hardware Faults at Fleet Scale
7.5 Check-and-Verify Defenses Against Silent Data Corruption
7.6 Software Faults
7.7 The Cost of Failure
7.8 Checkpointing and Progress Preservation
7.9 Failure Detection and Recovery
7.10 Elastic Recovery
7.11 Serving Fault Tolerance
7.12 Graceful Degradation
7.13 Distributed Debugging and Observability
7.14 Fault Injection Tools and Frameworks
7.15 Case Studies
7.16 Fallacies and Pitfalls
7.17 Summary

Purpose

Why does scale transform hardware failure from rare exception to routine condition that systems must absorb continuously?

A single accelerator can run for years between hardware failures. A thousand accelerators turn that same component risk into failures every couple of days, and ten thousand see failures every few hours, while shared power and network faults can interrupt many workers at once. Individual devices need not become less reliable for the fleet to experience failures routinely. A program on one machine can often stop and restart when a component fails; a distributed training job must also recover agreement among surviving workers about which progress remains valid. Restarting quickly is of little value if the restored state silently changes the computation. Preserving more state reduces the work at risk, but consumes storage bandwidth and memory while the job is healthy. Waiting longer between saves reduces that overhead while increasing the progress a failure can erase. Serving imposes a different consequence: users may abandon a request before recovery completes, and replacing a worker may leave the request's context unavailable. A response that arrives on time can still be unacceptable if recovery has discarded information needed to produce it. Engineers must therefore bound lost work and recovery delay while preserving the validity of training and the usefulness of service. The goal is continued useful execution despite repeated failures, with the costs and limits of recovery made explicit. In C^3 terms, fault tolerance spends compute and communication on coordination so that component failures do not routinely become failures of the entire workload.

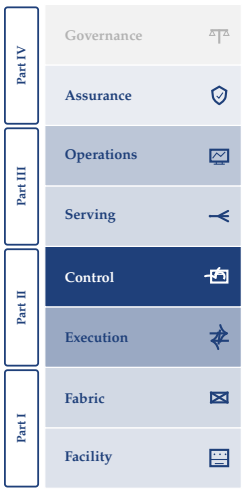

Learning Objectives

- Calculate mean time between fleet failures from component rates and estimate interruption frequency for large training clusters
- Classify hardware, software, silent-data-corruption, and slow-but-alive failures by detectability, blast radius, and recovery path
- Select checkpoint intervals that balance write overhead against acceptable lost work and recovery delay
- Design distributed checkpointing and elastic recovery plans for multi-terabyte model state across accelerator fleets
- Design serving redundancy, failover, and state-recovery policies under latency budgets and partial failures
- Select graceful degradation strategies that preserve useful service when models, features, regions, or capacity fail
- Evaluate detection timeouts, restart strategies, and fault-injection evidence against a per-failure recovery-time budget

7.1 Failure Analysis at Scale

Imagine a 10,000-GPU cluster midway through a three-month training run for a new foundation model. The parallel training plan and the communication layer beneath it have done their job, and thousands of devices exchange gradients through AllReduce, AllGather, and AllToAll as if they were one machine. Then the arithmetic of scale catches up. A GPU fails every few hours, and if the system cannot absorb that ordinary physical event, synchronized training halts while compute worth about $0.5M a day sits idle. **Fault tolerance** is the system property that lets distributed execution continue making useful progress when components fail, degrade, or disappear. In the fleet stack of figure 1.4, fault tolerance occupies the Reliability & Control row of the Distribution band, where it protects the progress that parallelism and communication produce.

That fragility is the direct consequence of successful synchronization. Every collective in chapter 5 assumed that all ranks in its process group arrive, and each dependent consumer waits for every contribution. One stalled device can therefore stop progress in its synchronous training group, even while unrelated jobs continue elsewhere in the fleet. Fault-tolerant ML systems preserve progress by detecting failure, checkpointing recoverable state, restarting or elastically reshaping jobs, and treating hardware churn as normal execution rather than an exceptional path. That reliability envelope covers both halves of the fleet's work. Training needs recovery that preserves weeks of accumulated progress, while serving needs redundancy, failover, and graceful degradation that preserve a latency contract, and both draw their spare capacity from a shared fleet pool.

↪ How the scheduler holds spare capacity for failed jobs and prepositions replacement nodes while a missed heartbeat is confirmed is developed in section 8.2.2.

Not every failure announces itself. A component can stop outright, slow down while still answering liveness checks, or keep running while it computes wrong values. The quiet cases pair low detection rates with a high blast radius, because the system keeps running on state it should no longer trust. The recovery target here is therefore trustworthy execution: preserving valid training progress and restoring a service after its infrastructure or software fails. A restored model can still make poor predictions on changed inputs even when every component executes correctly. Evaluating that behavioral robustness requires a separate contract for model quality under distribution shifts and adversarial inputs, developed in chapter 14. The distinction lets us use model-quality signals to diagnose a fault without treating every quality regression as evidence of failed hardware.

Because failures cannot be eliminated at this scale, fault-tolerant systems verify completion rather than assume it, treat failure as a normal code path, exercise recovery continuously, and account for partial failures, in which some components fail while others keep running, a case naive error handling may not anticipate. These techniques draw on decades of distributed systems research, but ML workloads change the economics. Training checkpoints are large but can be sized from the model and optimizer state. Some optimization algorithms also tolerate bounded numerical perturbations without requiring bitwise replay. That tolerance must be established for the training method; stochasticity does not make arbitrary corruption or inconsistent recovery safe. A recovery

design can exploit approximate execution only while preserving its declared convergence and quality requirements. Making failure a routine part of operation starts with knowing how often failures arrive, and the mathematics that follows makes the shift from *rare exception* to *routine occurrence* precise.

7.2 The Mathematics of Inevitable Failure

A single data-center GPU can run for years without a hardware fault, yet a synchronous job spanning ten thousand of those same GPUs is interrupted every few hours. System reliability engineering explains why individual device longevity erodes so quickly when thousands of components must stay healthy at the same time (Birolini 2017). Each component has a failure rate λ,[1] measured in failures per unit time. When λ is constant, the probability that a single component survives without failure until time t follows an exponential distribution,[2] as in equation 7.1:

$$R_{\text{single}}(t) = e^{-\lambda t} \tag{7.1}$$

The expected time to failure of this component equals $1/\lambda$. In this book's convention, mean time to failure (MTTF) names a single component's expected time to its first failure. MTBF names the repairable system's expected operating time between successive failures, excluding recovery downtime. Under the constant-rate exponential model used here the two coincide numerically, and the fault-tolerance model uses MTTF for per-component constants and MTBF for composed system rates. Data center GPUs often use planning MTTF values in the tens of thousands of hours, with field behavior depending on operating conditions, cooling effectiveness, manufacturing variation, and workload stress.[3] Section E.1.2 catalogs the reference per-component FIT and MTTF values (GPU, high-bandwidth memory (HBM), network interface card (NIC), power supply, PCIe switch, optical cable, ToR switch) that anchor these λ and FIT figures. A reader can substitute the constant for any device and reproduce the rates used throughout the fault-tolerance analysis.

When multiple independent components operate in a system where any single failure causes system failure, equation 7.2 formalizes how system reliability becomes the product of individual component reliabilities:

$$R_{\text{system}}(t) = \prod_{i=1}^{N} R_{\text{single}_i}(t) = \prod_{i=1}^{N} e^{-\lambda_i t} \tag{7.2}$$

For N identical components with individual failure rate λ, equation 7.3 gives the identical-component simplification:

$$R_{\text{system}}(t) = e^{-N\lambda t} \tag{7.3}$$

The system failure rate becomes $N\lambda$, and equation 7.4 expresses how the system MTBF scales inversely with component count. This inverse scaling explains why per-device reliability that looks excellent in isolation erodes at cluster scale:

$$\text{MTBF}_{\text{system}} = \frac{1}{N\lambda} = \frac{\text{MTTF}_{\text{component}}}{N} \tag{7.4}$$

The consequence is a gap between failure cadence and job length. At 10,000 GPUs, the GPU-only calculation gives an expected interval of 5 hours, so a training run measured in weeks accumulates dozens of expected failures and one measured in months accumulates hundreds. These are planning averages, not a schedule of incidents; a shorter job can finish without failure even when its expected failure count exceeds one. Once the expected time between failures falls below the length of the job, fault tolerance becomes part of the execution budget for long-running fleet work. Applied to the landmark training runs of section 1.1, whose accelerator counts grew by about four orders of magnitude between 2012 and 2024, the same per-device planning rate puts the largest of them, with more than ten thousand GPUs, below a 10-hour system MTBF. Recovery therefore becomes an active scheduling constraint rather than a reliability feature.

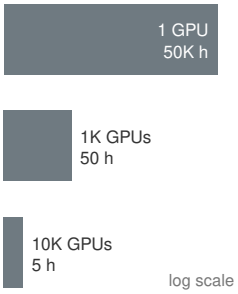

Fleet mean time between failures (MTBF) collapses from 50,000 h to 5 h as GPUs multiply.

1 **Failure rate (λ)**: Expressed in FITs (Failures In Time), where 1 FIT equals one failure per billion device-hours. A data center GPU with a 50,000-hour mean time to failure has FIT $= 20{,}000$, corresponding to $\lambda_{\text{hour}} = 2.0 \times 10^{-5}$ failures/hour. That seems negligible for one device but becomes dominant at fleet scale. A cluster with 10,000 GPUs accumulates 200,000,000 FITs from GPUs alone, translating to an expected GPU failure every 5 hours.

2 **Poisson process**: A statistical model for events occurring independently at a constant average rate. In this model, event counts follow a Poisson distribution and waiting times follow an exponential distribution, giving the survival function $R_{\text{single}}(t) = e^{-\lambda t}$. With independent identical component hazards, the probability of at least one failure in a 10,000-component cluster is $1 - e^{-N\lambda t}$. The aggregate hazard $N\lambda$ scales linearly with component count; the failure probability is nonlinear in both N and t, approaching $N\lambda t$ only when that product is small.

↪ The independence, exponential-lifetime, and Poisson results this failure arithmetic rests on are developed in section E.1.1.

7.2.1 Composing component reliability

The GPU-only calculation of equation 7.4 supplies the scale intuition; a real training system then composes GPUs, HBM, links, power supplies, storage paths, and redundancy into one failure budget. Component rates are usually quoted in the FIT units introduced with the failure rate, so the 50,000-hour GPU contributes 20,000 FIT to whatever system contains it. The new question is how different component rates compose, and the composition rule depends on the redundancy structure and the independence assumptions. For series systems (for example, a node where all 8 GPUs must work), failure of any component fails the whole. With independent constant-rate failures, the reliabilities multiply as in equation 7.2, so the rates add even when the components differ:

$$\text{MTBF}_{\text{system}} = \frac{1}{\sum \lambda_i}$$

Under this additive failure-rate law, the system failure rate grows linearly with component count while system MTBF falls in inverse proportion. Every non-redundant component on the critical path adds its FIT contribution directly to the denominator.

The same product governs instantaneous availability, a different quantity from the survival function $R_{\text{single}}(t)$. Availability A is the probability that a component is up at a randomly chosen instant, counting repairs, whereas $R_{\text{single}}(t)$ is the probability that it has not yet failed by time t. If each GPU is independently up with probability A_{GPU}, all N are up together with probability $(A_{\text{GPU}})^N$, which is the reliability gap of equation 1.2, in which high per-device availability still yields low availability for a job that needs every device. For the chapter's 10,000-GPU cluster, a per-GPU availability of 99.99 percent (only 52.6 minutes of downtime per year) leaves about a 37 percent chance that every GPU is up at the same instant. This compounding erodes even highly reliable hardware once N reaches fleet scale, so software must handle failures automatically.

> **Systems Perspective 7.1: Scale transforms failure**
>
> A cluster of 10,000 GPUs turns a per-GPU MTTF of 50,000 hours into a system MTBF of 5 hours. For a repairable system, steady-state availability is $A = \text{MTBF}/(\text{MTBF} + \text{MTTR})$, where mean time to recovery (MTTR) is the recovery interval. Cutting MTTR from 30 minutes to 5 minutes raises availability from 90.91 percent to 98.36 percent. This availability treats the whole job as one repairable system that resumes after each recovery. It therefore measures a different quantity from the instantaneous probability, computed above from a per-GPU availability, that every GPU is up at once. The stationary two-state model assumes unconstrained repair capacity and complete outage; correlated incidents, repair contention, and degraded operation require direct measurement (Birolini 2017). Systems must engineer recovery, not merely improve components.

Engineering faster recovery is one response to that arithmetic; redundancy is the other. For parallel systems providing redundancy, suppose that any surviving component can satisfy the required service. With independent failures and no repair during the observation interval, the system survives until every component has failed:

$$R_{\text{system}}(t) = 1 - \prod_{i=1}^{N_{\text{red}}} (1 - R_{\text{single}_i}(t))$$

Series composition is the training job's problem, because a synchronous step needs every participant. Parallel composition is the serving fleet's lever, and it returns when interchangeable replicas protect a latency contract. Both rules, however, describe whatever the model calls a component, and the fault rate assigned to that component depends on which protections sit inside it.

[3] **GPU MTTF variation**: This chapter uses 50,000 hours as a planning constant for current data center GPUs (A100 and H100 class), not as a guarantee for every deployed GPU. Published fleet studies from Google TPUv4 pods and Meta GPU clusters report that field reliability depends on topology, cooling, workload, and operational practice (Zu et al. 2024; Kokolis et al. 2025; Dubey et al. 2024). The planning constant therefore anchors the arithmetic, while operators should substitute measured fleet rates when they have them.

Lighthouse 7.1: Archetype A (GPT-4/Llama-3): Memory soft-error scale

Composition makes the distinction between raw faults and job interruptions essential. Memory soft-error rates are quoted per megabit of storage rather than per device, so the exposure grows with memory capacity. As an illustrative raw-bit-error budget, assume 250 FIT per megabit across 1,024 GPUs with 80 GB each, the size of the Archetype A reference cluster. The resulting interval between raw memory errors is roughly 20 seconds. This is not a measured job-failure interval. Error-correcting code (ECC) protection can correct covered error patterns before they affect computation, while detected uncorrectable errors and escaped corruption require different recovery paths. At that raw rate, unprotected memory in the reference cluster would corrupt training state many times per hour, so ECC is a precondition for large-scale training rather than an option. A fleet reliability budget therefore needs the residual job-interrupting and silent-corruption rates under its actual protection scheme; raw bit-error rates alone cannot establish whether memory, logic, or interconnect dominates.

These raw bit-error rates justify ECC, but they are not the interruption rates that set checkpoint cadence; the worked cascade that follows uses device MTTF values that already include the effect of that protection. The cascade also treats its GPU, PCIe, power, and network contributions as separate categories, so a measured whole-node rate that already includes those paths cannot be added to them again.

7.2.2 Cluster MTBF calculation

Consider a training cluster designed for large language model development, roughly ten times the 1,024-GPU Archetype A reference cluster, with the following specifications:

- 10,000 NVIDIA H100 GPUs
- Individual GPU MTTF: 50,000 hours
- Each GPU connected to host via PCIe (MTTF: 200,000 hours)
- Each node contains 8 GPUs fed by 6 power supply units (MTTF: 100,000 hours each) in 4+2 redundancy, so the node keeps full performance with any 2 of them failed
- Network infrastructure per node (NIC, cables): MTTF 150,000 hours

Each GPU operates within a failure domain that includes the GPU itself, its PCIe connection, and a proportional share of the node's shared power and network infrastructure. Dividing a node-wide subsystem rate by the 8 GPUs in the node prorates one shared failure source across those GPUs, so that multiplying the per-GPU total by all 10,000 GPUs in the final step counts each node's shared hardware once per node rather than eight times. Power enters the sum only through supplies whose single failure stops the node, counted as n_{psu} in the power equation. With the node's supplies in redundant sets, no single unit qualifies, so the power term below is zero in this first-order calculation. The approximation omits the residual risk that additional independent failures exhaust redundancy before repair. A fault upstream of the chassis, such as the loss of the circuits feeding it, can also take down the whole node. That risk belongs to the correlated failures treated later in this chapter. The four subsystem rates follow from their MTTF values:

$$\lambda_{\text{GPU}} = \frac{1}{50{,}000} = 2.0 \times 10^{-5} \text{ failures/hour}$$

$$\lambda_{\text{PCIe}} = \frac{1}{200{,}000} = 0.5 \times 10^{-5} \text{ failures/hour}$$

$$\lambda_{\text{power/GPU}} = \frac{n_{\text{psu}}}{8} \times \frac{1}{100{,}000} = \frac{0}{8} \times \frac{1}{100{,}000} = 0 \text{ failures/hour}$$

$$\lambda_{\text{network/GPU}} = \frac{1}{8} \times \frac{1}{150{,}000} = 0.083 \times 10^{-5} \text{ failures/hour}$$

Because the subsystems sit in series, their rates add to a per-GPU total:

$$\lambda_{\text{total/GPU}} = (2.0 + 0.5 + 0.083) \times 10^{-5} = 2.583 \times 10^{-5} \text{ failures/hour}$$

Multiplying that total by the cluster size gives the system failure rate, whose reciprocal is the cluster MTBF:

$$\lambda_{\text{system}} = 10{,}000 \times 2.583 \times 10^{-5} = 0.2583 \text{ failures/hour}$$

$$\text{MTBF}_{\text{system}} = \frac{1}{0.2583} = 3.87 \text{ hours}$$

Under the model, the cluster experiences a failure every 3.9 hours on average. The expected failure cadence is 6.2 failures/day; a training run lasting one week has an expected count of 43 failures. Actual incident times vary, but the expected loss is large enough that recovery must be part of normal execution planning.

↪ The step-by-step composition of per-component failure rates into a fleet-wide system MTBF, reusable for any node design or cluster size, is developed in section E.1.3.

The cluster MTBF scaling in table 7.1 isolates the GPU-only baseline across cluster sizes. The full-system calculation is lower once PCIe and network failure domains are included. The same 10,000-GPU cluster therefore carries three MTBF values in this chapter. Each later calculation names the one it uses: the GPU-only 5 hours of this table, the full-system 3.9 hours of the cascade above, and the two-to-three-hour cadence that somewhat larger production runs report below. The full-system MTBF is the default input; a calculation that uses the GPU-only value says so, usually because it compares cluster sizes on the common basis of table 7.1, and its failure count is then a lower bound.

Two running scenarios carry the rest of the chapter. The 1,024-GPU Archetype A reference cluster sizes per-node state, restore paths, and the elastic and straggler examples, while the 10,000-GPU cluster of the cascade sets failure cadence and checkpoint intervals.

Table 7.1: Cluster MTBF Scaling: GPU-only mean time between failures falls in inverse proportion to cluster size, transforming failures from rare events to continuous operating conditions at scale. Additional host, power, storage, and network failure domains reduce system-level MTBF further.

Cluster size (GPUs)	Individual GPU MTTF	Cluster MTBF	Expected failures per day
8	50,000 hours	6,250 hours (260 days)	0.004
64	50,000 hours	781 hours (33 days)	0.03
512	50,000 hours	98 hours (4 days)	0.25
1,000	50,000 hours	50 hours (2 days)	0.48
4,000	50,000 hours	12 hours	1.9
10,000	50,000 hours	5 hours	4.8
25,000	50,000 hours	2 hours	12

Production measurements confirm the $1/N_{\text{GPU}}$ trend of table 7.1 and show that the GPU-only line is optimistic. Plotted on log-log axes, measured MTTF values and fitted projections from Meta's Research SuperCluster (RSC) fall roughly along the slope of the $50{,}000\,\text{h}/N_{\text{GPU}}$ line across more than three orders of magnitude in job size, though less steeply between the largest measured job and the first projection, and sit below it. The gap arises because a job can be interrupted by far more than its GPUs (figure 7.1). The Kokolis et al. (2025) study of the RSC reported measured MTTF values for observed RSC job sizes and projected a 1.8-hour MTTF for 16,384-GPU jobs. That projection lands next to the Llama 3 run previewed in section 1.1 (war story 7.2 follows its recovery). That run recorded 419 unexpected interruptions across a 54-day snapshot on 16,384 H100 GPUs, which gives an observed MTTF of about 3.1 hours. Unlike the RSC values, that measurement sits close to the GPU-only line at its job size rather than below it.

Together, the measured Llama run and Kokolis projection put 16,384-GPU training on a failure cadence of roughly two to three hours, bracketed by the 1.8-hour projection and the 3.1-hour measurement. This evidence turns the idealized $1/N_{\text{GPU}}$ relation into an engineering constraint that determines checkpoint frequency, recovery architecture, and infrastructure investment.

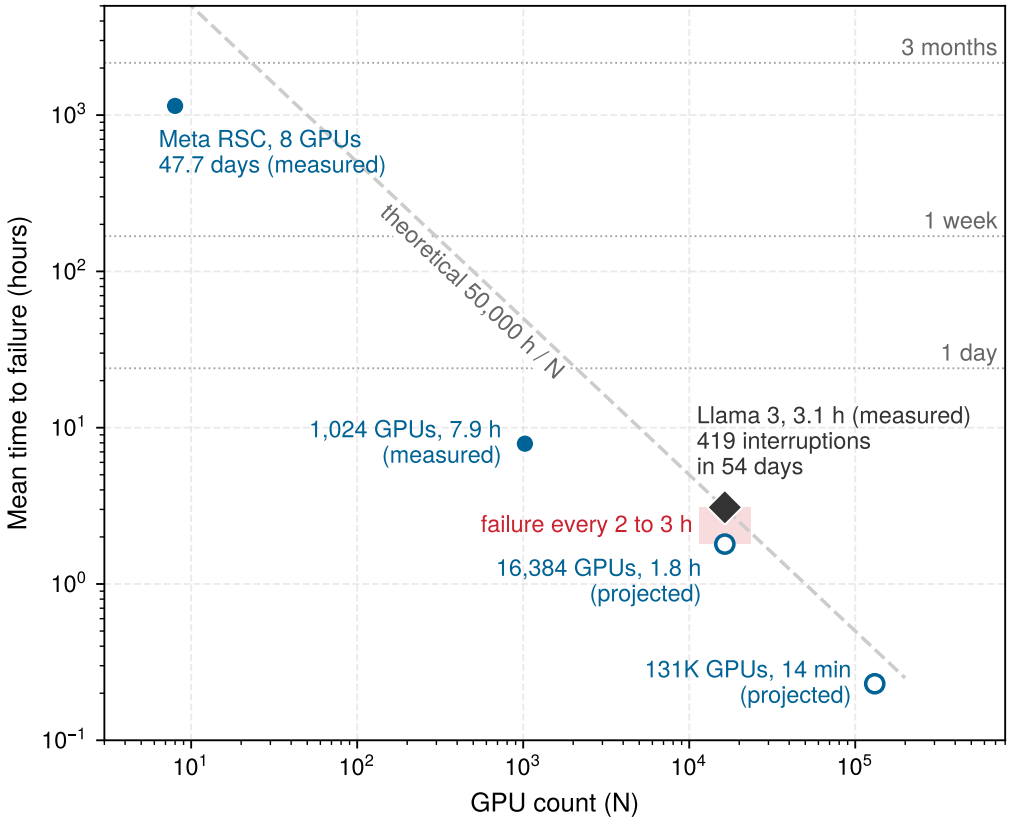

Figure 7.1: Published Failure Rates from Production Clusters: Measured MTTF values and fitted projections from Meta's RSC clusters (Kokolis et al. 2025) plotted with Llama 3 training measurements (Dubey et al. 2024) against the theoretical 1/N curve. At 16,384 GPUs, the Kokolis projection and the Llama 3 empirical measurement bracket a failure cadence of roughly two to three hours (the 1.8-hour projection and the 3.1-hour measurement), confirming that fault tolerance is a mathematical necessity, not an edge case. Data sources: Kokolis et al., arXiv:2410.21680; Meta, arXiv:2407.21783.

7.3 Failure Taxonomy

A network partition that clears in seconds and a permanent GPU failure both interrupt a training step, yet retrying the first succeeds while retrying the second wastes cluster time on dead silicon. Silent memory corruption that produces wrong gradients demands a different defense again from a node crash that stops responding altogether. While the MTBF calculations quantify *how often* interruptions arrive, which sets checkpoint intervals and spare-pool sizes, choosing the right recovery action requires distinguishing *what kind* of failure has occurred. In standard reliability terms, a fault is the underlying defect or disturbance, an error is the corrupted internal state it produces, and a failure is the externally visible misbehavior that the training or serving system must handle (Avizienis et al. 2004).

No single label captures all of these differences. Reliability decisions require three orthogonal axes (Constantinescu 2008; Lamport et al. 1982): persistence (transient, intermittent, or permanent), manifestation (fail-stop, degraded, or value-corrupting), and failure scope (one device, one node, a rack or power domain, or the whole job). Two manifestations need names because liveness checks cannot see them. A **gray failure** is a component that stays alive but runs slowly or erratically, so it keeps answering the liveness checks that would flag a crash while every peer that waits on it slows down. **Silent data corruption (SDC)** goes further, producing wrong values with no error at all.

Because the three axes are independent, a single fault sits on all three at once. A transient disturbance can manifest as either a detected crash or silent corruption, while a permanent switch failure can remain confined to one rack or partition an entire job. Each axis governs a different operational

choice. Persistence determines whether a fault clears on its own or requires hardware quarantine and replacement (section 7.4.1). Manifestation separates crashed ranks from slow stragglers (section 7.9.5) and from Byzantine value corruption, while scope determines when a single shared hazard can defeat naive redundancy. The subsections that follow take one case from each axis: transient failures for persistence, fail-stop and Byzantine failures for manifestation, and correlated failures for scope.

Production logs show which of these failures a large run actually meets. During a 54-day period of Llama 3 405B pretraining on 16,384 H100 GPUs, Meta attributed each of 419 unexpected interruptions to a root-cause component (Dubey et al. 2024). Grouped by cause in figure 7.2, GPU components account for 268 of them, about 64 percent, led by 148 faulty GPUs and 72 HBM memory failures. Software caused 56, of which 54 were bugs and 2 were failed dependencies. Network switches and cables caused 35, unplanned host maintenance 32, and the remaining host hardware, from NICs to power supplies, 21. The last 7 were NCCL watchdog timeouts whose underlying cause was never found. Only 6 interruptions were attributed to silent data corruption, and that count includes only corruption that was caught. Most fault-tolerance effort therefore goes to frequent, detectable accelerator failures, while the rare silent ones demand a different kind of defense.

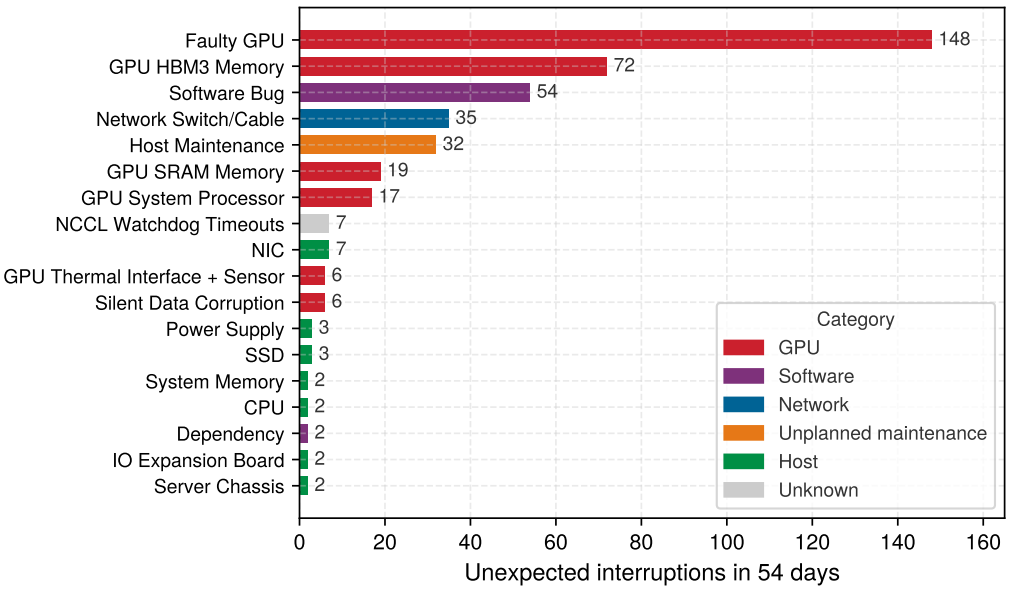

Figure 7.2: Root Causes of Training Interruptions: The 419 unexpected interruptions recorded over 54 days of Llama 3 405B pretraining on 16,384 H100 GPUs, by root-cause component and colored by category. GPU components dominate; silent data corruption appears only where it was detected. Counts are from Table 5 of Dubey et al. (2024); one percentage in that table (Faulty GPU) disagrees with its count, so shares here are computed from the counts.

7.3.1 Transient failures

A transient fault clears without physical repair, but the state it corrupts can persist after the disturbance ends, forcing recovery to retry or validate affected work before relying on that state. In large-scale training systems, temporary disruptions range from network packet loss, where dropped packets are recovered by transport retransmission, to memory bit flips triggered by cosmic-ray-induced single-event upsets.[4] Timing disturbances can be equally disruptive without corrupting bits directly. Under temperature spikes, thermal throttling reduces compute throughput without halting execution. Software timeouts fire when transient contention delays a collective beyond its allotted window.

Transient faults are particularly difficult to diagnose in ML training when they leave no explicit error signal. If hardware protection and the computation itself do not mask a memory bit flip during gradient computation, the corrupted value propagates into later updates while training continues. Large-scale SDC studies document such escaped corruption in production processors, and DRAM

[4] **Single-event upset (SEU)**: From particle physics, where "upset" denotes a nondestructive state change. Cosmic rays and alpha particles from chip packaging can flip bits in memory and logic; the observed rate depends on altitude, process technology, packaging materials, shielding, and workload (Ziegler et al. 1996; Baumann 2005). ECC memory corrects single-bit errors but cannot handle every multi-bit upset in the same word, leaving a residual silent-corruption risk that compounds across the terabytes of state in large-scale ML training.

field studies confirm that memory errors occur frequently enough to require fleet-level monitoring (Dixit et al. 2021; Sridharan et al. 2015; Schroeder et al. 2009). In ML training, that undetected corruption surfaces as degraded gradients, anomalous loss curves, or rollback-worthy checkpoints rather than an explicit crash.

How the system responds to a transient failure therefore depends on whether the fault announces itself. Errors that raise explicit exceptions can be handled through bounded retry logic, whereas silent corruption requires end-to-end validation. Checksums catch altered bytes in storage or transport, cross-rank digests, which compare a short hash of each rank's copy, expose disagreement across a supposedly identical reduced buffer, and statistical or replay checks test whether the computation itself remains plausible. Because a rank that computes the wrong local gradient and then checksums those wrong bytes will pass an ordinary transport check, only independent recomputation, invariants, or model-level evidence can expose the error.

7.3.2 Fail-stop failures

In synchronous distributed training, a single participant that goes silent stalls the entire group's next update. A conventional AllReduce cannot complete until every rank supplies its input, so one dead GPU in a 10,000-GPU job can leave the other 9,999 accelerators waiting at the barrier until recovery begins. Smaller communication groups limit the immediate dependency, although pipeline or data-parallel stages soon propagate the stall across the rest of the job. Such **fail-stop failures** occur when a component ceases operation outright and stops responding to requests. A GPU hardware failure renders an individual accelerator unresponsive, and a host node crash terminates all training processes on the machine. A network partition severs communication between a healthy node and the cluster, and a storage failure blocks workers from reading or writing checkpoint shards.

A stopped component gives a clearer symptom than one that keeps returning plausible but incorrect values. Yet a stopped rank announces itself only through silence, so detection rests on timeouts that must separate an outright crash from a slow or isolated host, such as a GPU still writing a large checkpoint while the collective waits. Once a timeout flags the missing participant, the recovery system replaces or restarts it and restores the most recent valid checkpoint; section 7.9.1 sets that timeout threshold once the full recovery-time budget is in view.

7.3.3 Byzantine failures

Where a fail-stop component ceases to participate, a value-corrupting component can keep running while returning incorrect results. A GPU may return wrong gradients without raising an error, a link may deliver a corrupted packet whose error pattern its cyclic redundancy check (CRC) happens not to cover, or a worker may compute different results for identical inputs and state. Byzantine failures allow arbitrary behavior, including inconsistent messages to different participants (Lamport et al. 1982). ML systems can encounter this broader failure model through silent data corruption or adversarial corruption. Numerical instability and determinism violations also require validation when they break the execution contract, but ordinary rounding differences or intentionally stochastic outputs are not themselves evidence of a faulty worker. The recovery concern is a participant whose accepted contribution can corrupt shared state without a visible crash.

The two models leave different evidence (figure 7.3). A fail-stop worker such as W2 falls silent, and its missing heartbeats, the periodic liveness messages each worker sends, make it suspect. A Byzantine worker such as W6 keeps sending messages on schedule, so it passes every liveness check while the values it sends are wrong.

In an ML fleet, the most common source of value corruption is not an adversary but hardware that changes a value without raising an error, the silent data corruption already met among transient failures. Its physical causes, the training signals that expose it, and the verification and hot-spare defenses that contain it are developed together in section 7.5, after the hardware faults that produce it. Whatever its source, a Byzantine contribution violates a contract that liveness checks never test. Ordinary data-parallel workers process different batches and need not produce identical local gradients; the required contract is that each contribution correctly represents its assigned computation. A single Byzantine worker can violate that contract and corrupt the averaged gradient, potentially causing training to diverge or converge to a poor solution. Section 13.6.1 examines the same weakness when the corrupting worker is a deliberate attacker.

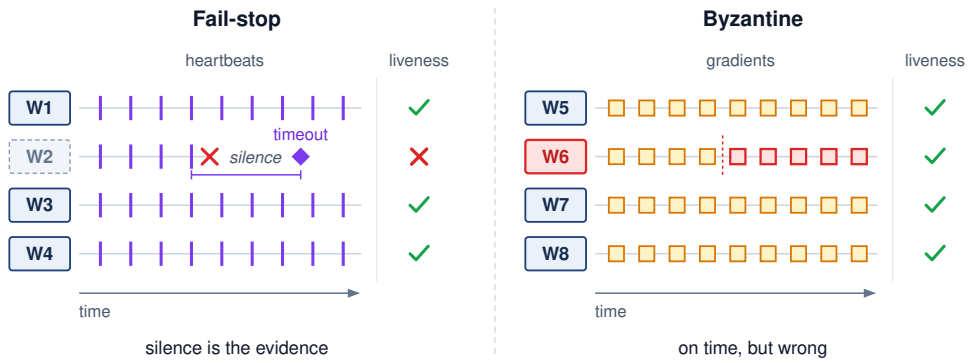

Figure 7.3: Fail-Stop vs. Byzantine Failures: In the fail-stop model (left), a failed worker simply ceases to send messages, which a timeout flags as suspect without distinguishing a crash from delay or isolation. In the Byzantine model (right), a failed worker continues to participate but sends incorrect data (for example, corrupted gradients reported as valid), which can poison the global model state if not detected by redundant validation.

Redundant computation supplies one form of evidence against Byzantine failures. Workers repeating a computation from matching inputs, state, and random-number streams can compare results under a stated numerical tolerance. Statistical outlier detection and invariants provide additional evidence without fully repeating the computation. Each method has a coverage limit. Agreement can miss shared faults, while a statistical anomaly can arise from legitimate data variation. Detection overhead must therefore be evaluated together with the faults the method can expose.

Byzantine-resilient distributed training algorithms exist but impose significant overhead. Krum (Blanchard et al. 2017) keeps the single worker gradient whose nearest neighbors lie closest to it and discards the rest, while coordinate-wise trimmed mean (Yin et al. 2018) drops the largest and smallest values in each coordinate before averaging what remains. Both compute aggregates that stay robust to a bounded number of Byzantine workers, but they require more communication and computation than simple averaging. Without such protection, corrupted gradients can push the optimizer toward an unreliable model state while the training job appears healthy. Reliability must therefore protect semantic correctness rather than process liveness alone.[5]

7.3.4 Correlated failures

The reliability calculations in section 7.2 assume independent failures, yet real systems depart sharply from this assumption when common physical or operational dependencies take down multiple components at once. Within a single node, a power fault that exhausts the chassis's redundant supplies, such as the loss of its feed circuits, removes every GPU in the chassis. Across a cluster, a top-of-rack network switch failure isolates all attached nodes, while a cooling system failure can trigger thermal shutdowns down an entire row of racks. Shared software and operational dependencies widen the blast radius further. A software bug in an updated CUDA driver can crash processes across nodes whenever they exercise the affected code path, and operator error during cluster reconfiguration can render an entire fleet unreachable.

Correlated failures violate the independence assumption underlying equation 7.3. A model that counts only independent device failures omits shared hazards and their potentially larger blast radius; incident frequency and the number of affected devices must be modeled separately. Correlated failures can also defeat redundancy strategies. Three replicas of a model provide no protection against a shared power failure if all three run on the same power domain.

Defending against correlated failures requires classifying each failure by its scope, the failure domain it takes down, and ensuring that redundancy spans independent domains. Common failure domains in ML infrastructure range from a single GPU to an entire data center region, and each calls for its own mitigation (table 7.2). These domains nest, so a physical disruption scales upward from one eight-GPU node to a rack and then to a multi-rack power and network envelope (figure 7.4). The software-version domain in the table cuts across that physical nesting, because every node

[5] **Byzantine-resilient ML**: Named after Lamport's 1982 "Byzantine Generals Problem," these algorithms (Krum, trimmed mean, and signSGD, which aggregates only the signs of gradient coordinates by majority vote) tolerate a bounded fraction of corrupted workers, with the exact bound depending on the rule and its assumptions. Krum's original analysis requires $2f + 2 < n$ for f Byzantine workers among n, while coordinate-wise median and trimmed mean rest on different robustness assumptions, so the tolerated fraction depends on the aggregation rule, dimensionality, and statistical model. For n workers, Krum requires pairwise distance computations over worker gradients, so overhead grows quadratically with worker count and directly competes with the throughput gains of data parallelism (Blanchard et al. 2017).

running the same version shares its bugs wherever it sits. Its mitigation is a staged rollout, which first deploys a new version to a small canary slice of the fleet and widens the deployment only while that slice stays healthy.

↪ How a staged rollout limits the share of the fleet a new version reaches while evidence of its health accumulates is developed in section 12.5.5.

Table 7.2: Failure Domains in ML Infrastructure: Understanding failure domain boundaries enables placement of redundant components across independent domains, preventing correlated failures from defeating redundancy strategies.

Failure domain	Impact scope	Typical recovery time	Mitigation strategy
Single GPU	1 GPU	Seconds (spare)	Hot spare, elastic training
Node (power/OS)	8 GPUs	Minutes	Checkpoint, node replacement
Rack (ToR switch)	32–64 GPUs	Minutes to hours	Cross-rack redundancy
Power domain	100–500 GPUs	Hours	Multiple power feeds
Data center region	All GPUs	Hours to days	Geographic distribution
Software version	All GPUs running version	Minutes (rollback)	Staged rollouts

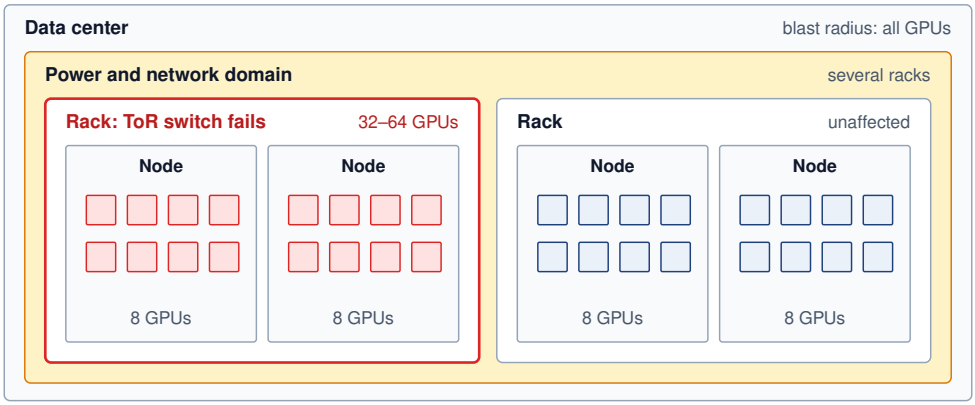

Figure 7.4: Hierarchy of Failure Domains: Failure domains are often nested or overlapping. A GPU failure affects one device. A node failure affects 8 GPUs. A rack switch failure affects 32–64 GPUs (each rack is drawn with only two of its four to eight nodes). A power distribution unit (PDU) failure may affect multiple racks. The red rack shows one such correlated failure: a single top-of-rack switch fault removes every GPU it contains, while the neighboring rack is untouched, which is why replicas must span independent domains (for example, different racks or rows).

The containment hierarchy in table 7.2 and figure 7.4 explains why replica placement must match the fault being tolerated. Placing replicas within the same domain provides zero protection against that domain's failure mode. The highlighted rack in figure 7.4 contains nodes that share a top-of-rack (ToR) switch, so that switch's loss can disconnect every attached node from the job. A neighboring rack with a separate switch can remain reachable during that event even if both racks share a power domain, yet that same placement does not protect against loss of their shared power source. Redundancy must therefore cross the relevant failure boundary, not merely occupy different devices.

↩ The failure domain and its physical nesting from node to rack to building are established in section 2.5.5.

Checkpoint 7.1: Mapping failure domains

Verify your understanding of how failure domains nest and their operational impact:

- ☐ **Blast radius**: calculate how many 8-GPU nodes a rack switch failure typically affects in a standard data center configuration.
- ☐ **Power-domain correlation**: explain why placing all 3 model replicas in the same power domain fails to provide protection against a power outage confined to that domain.

- ☐ **Staged rollouts**: identify the failure domain addressed by staged rollouts rather than hardware redundancy.
- ☐ **Nested domains**: explain hierarchical containment and decide whether the rack failure domain still matters once the power domain containing that rack fails.

War Story 7.1: The supercomputer that swapped in spare cubes (2024)

Context: Google's Tensor Processing Unit (TPU) v4 machine-learning supercomputer used 4,096 TPU chips connected by a custom 3D torus inter-chip interconnect (ICI) fabric, and Google engineers published their operating experience with automatic fault resiliency and hardware recovery (Zu et al. 2024).

Mechanism: At 4,096-chip scale, machine and link failures could remove participants or paths required by a job's configured communication topology.

Impact: Losing a required participant can interrupt the synchronous job across the slice. Losing a link instead raises a routing question: whether the remaining fabric can still provide the required communication paths.

Fix: Google used optical circuit switches (OCS) to assemble replacement slices from spare healthy TPU cubes, the pod's building blocks of chips wired as a $4\times4\times4$ mesh, for machine outages, together with fault-tolerant routing for supported OCS outage patterns. The response depended on which resource failed.

Systems lesson: Fault tolerance is a topology-management problem. In tightly coupled accelerator fleets, recovery must preserve the communication graph, not merely replace an individual component; otherwise one local hardware fault can become a cluster-wide training interruption.

7.4 Hardware Faults at Fleet Scale

A cosmic ray that flips one bit in a GPU's high-bandwidth memory and a solder joint that cracks under thermal cycling are physical events, but the recovery system never observes them directly. It observes symptoms: a gradient spike with no process crash, a rank that disappears from the collective, or a node that fails only after it heats under sustained load. The failure taxonomy of section 7.3 already classifies such events by persistence, manifestation, and scope. For hardware, persistence carries most of the operational weight because it decides whether the fleet should retry, quarantine, or replace a component, while manifestation decides whether the fleet notices the fault at all.

ML workloads raise the stakes of each fault. A training step performs trillions of arithmetic operations on state that persists into later steps, so a wrong value can survive beyond the computation that produced it. The damage depends on the affected value and its later use. Under the IEEE 754 format,[6] a sign-bit flip inverts a weight, while an exponent-bit flip can move a small gradient by many orders of magnitude. Fault-injection studies show that network resilience varies with model, layer, and bit position, so a few faults in vulnerable locations can cost disproportionate accuracy (Reagen et al. 2018). A valid-looking corrupted value can enter learned state without triggering the process failure that would otherwise initiate recovery.

Fleet scale turns these rare device events into routine operating conditions. Smaller geometries and lower operating voltages can reduce the charge needed to disturb a stored value (Baumann 2005). Fleet studies at Google and Meta found that a small but persistent population of processors, on the order of a few per several thousand machines, computes wrong results without raising any error (Hochschild et al. 2021; Dixit et al. 2021). Studies of accelerator training fleets likewise conclude that rare hardware faults appear often enough that software-level detection and recovery become part of the training system design (He et al. 2023). ML system architects must therefore treat hardware as an unreliable substrate and budget for end-to-end integrity checks, replayed computation on high-value paths, and periodic consistency checks as system requirements rather than high-performance computing specialties.

[6] **IEEE 754**: The IEEE 754 floating-point standard defines the binary32 format with 1 sign bit, 8 exponent bits, and 23 fraction bits (IEEE Standards Association 2019). Because bit positions differ so much in consequence, resilience mechanisms often prioritize the sign and exponent bits rather than treating all bits as equally important.

↪ How an FP16 exponent-bit flip rescales a value by 2^{16} and can saturate an attention softmax is developed in section E.6.

7.4.1 Fault persistence and the fleet response

Each persistence class from the taxonomy maps to a different fleet action, from retry for a fault that leaves nothing to repair, through quarantine for one that returns under load, to replacement for one that repeats on every attempt (figure 7.5). A transient fault, as section 7.3.1 established, leaves no lasting defect from that event, so the fleet's question is only whether its effect escaped. A single-event upset, a voltage droop on an unstable power path (Reddi and Gupta 2013), a timing violation, or a soft error in combinational logic (Mukherjee et al. 2005) can corrupt a tensor, packet, or instruction. Detection and correction can prevent propagation; affected state may otherwise require rollback. During training, an undetected error in weights or gradients can compromise convergence and contaminate subsequent checkpoints (He et al. 2023). During stateless inference, an activation error can remain confined to one request (Mahmoud et al. 2020). If that activation enters retained state such as a KV cache, later outputs can also depend on it.

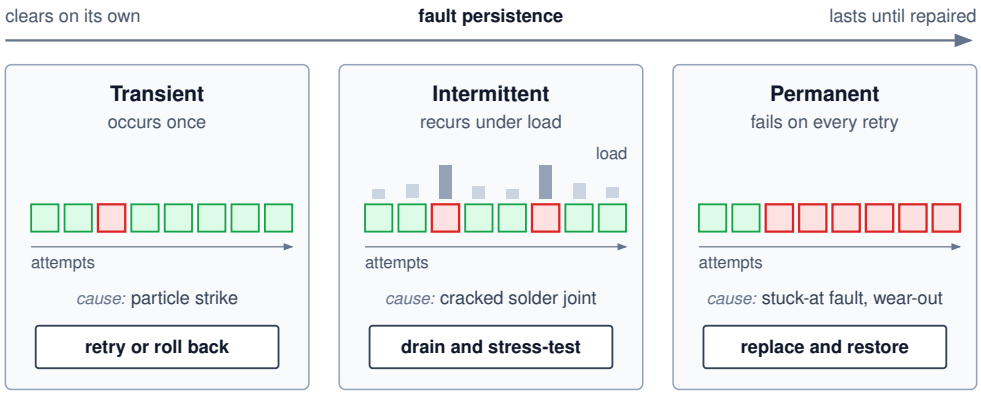

Figure 7.5: The ML Failure Spectrum: Hardware faults ordered by how long they persist. A transient fault leaves nothing to repair, an intermittent fault returns under the load that exposed it, and a permanent fault repeats on every retry. Persistence decides the fleet response, from retry through quarantine to replacement, and fault tolerance engineering must handle all three to keep a fleet making progress.

A permanent fault announces itself through repeatability. The same accelerator, memory cell, or link fails again after retry, often at the same address or in the same workload phase. Manufacturing defects and wear-out mechanisms such as electromigration[7] and oxide breakdown produce this signature, and the stuck-at fault, in which a signal remains fixed at 0 or 1 regardless of input, is its standard abstract model (Seong et al. 2010). Repeatability helps locate the defect but does not make the result trustworthy. A permanent datapath fault keeps producing wrong gradients until the device is drained from the cluster (He et al. 2023), and rerunning the computation on the same device reproduces the error instead of exposing it. The fleet must stop trusting the component, quarantine it, and restore from a checkpoint that predates its first bad write.

Intermittent faults create the most operational ambiguity because they produce evidence and then disappear. A node may pass a reboot test, rejoin the fleet, and fail again only when the next job drives the package into the same thermal, voltage, or communication regime. Cracked solder joints, aging packages, and contamination between contacts create these load-dependent conditions (Constantinescu 2008; Rashid et al. 2015), and reduced voltage margins cause timing errors that are hard to reproduce (Zhang et al. 2018). Sporadic errors of this kind can accumulate across iterations and degrade convergence without any explicit failure (Rashid et al. 2015). An intermittent node should therefore remain suspect until evidence clears it. The fleet drains it from long jobs, correlates its error logs, and readmits it only after it passes stress tests under the load that exposed the fault. Otherwise one marginal component can make validation and recovery untrustworthy for every job that lands on it. Across the three persistence classes, the evidence that exposes a fault also determines the fleet's response (table 7.3).

7 **Electromigration**: Gradual displacement of metal atoms in conductors by electron momentum transfer, first characterized in early electromigration studies and later modeled in Black's reliability equation (Black 1969). Mean time to failure decreases with higher current density and higher temperature, so sustained high-power accelerator workloads make thermal management a direct determinant of fleet lifespan.

Table 7.3: Hardware Fault Persistence and Fleet Response: The persistence class of a hardware fault determines the evidence that exposes it and the action the fleet takes. Manifestation is a separate axis: any of the three classes can surface as a crash or as silent corruption (Constantinescu 2008).

Persistence	Fleet signature	Detection evidence	Fleet response
Transient	One-time wrong value or lost message; retry succeeds	Corrected-error counters, checksum mismatches, gradient spikes	Correct or retry; roll back if durable state was contaminated
Intermittent	Recurring errors under specific load, heat, or voltage	Correlated error logs and repeated anomalies on one node	Drain and quarantine; readmit only after screening under load
Permanent	Reproducible failure at the same device, address, or link	Failed self-test or replay mismatch on the same device	Replace the component; restore from a checkpoint before the fault

7.4.2 Hardware lifecycle and the hazard rate

A newly unboxed accelerator rack is far more likely to fail in its first week of service than in its sixth month, and a multi-year-old rack that has endured thousands of thermal cycles becomes fragile again as packages and solder joints fatigue. While persistence classifies a fault after it strikes, the hazard function $h(t)$, the failure rate at a given age among components that have survived to that age, describes *when* over a component's operating lifespan faults are most likely to occur, partitioning hardware longevity into three reliability regimes (figure 7.6). The first phase, infant mortality, brings elevated failure rates from manufacturing defects, improper installation, and early-life failure of marginal components. For electronic components it typically lasts days to weeks, which is why burn-in testing[8] stresses hardware before deployment to precipitate those early failures prior to production use.

[8] **Burn-in testing**: A controlled screening procedure stresses components before deployment to expose early-life defects. Temperature, voltage, duration, and acceptance criteria depend on the component and its qualified test procedure; a generic semiconductor screening recipe is not an operating specification for an assembled accelerator. Screening reduces the defective population entering service but does not certify every later workload or fault mode.

After surviving infant mortality, components enter the useful-life phase, where failure rates remain relatively constant under the standard reliability-engineering model (Birolini 2017; Klutke et al. 2003). This phase spans the longest portion of a component's lifetime and is the regime where the exponential reliability model in equation 7.1 applies most accurately. For data center GPUs, the useful-life window is an operational planning concept shaped by refresh cycles, cooling, utilization, and observed fleet health rather than a fixed physical duration.

As components age, they enter the wear-out phase, where failure rates increase due to accumulated wear. For GPUs, wear mechanisms include electromigration in circuits, thermal cycling stress on solder joints, and degradation of thermal interface materials. The onset of wear-out depends heavily on operating conditions; components operated at high temperatures or with frequent thermal cycling enter wear-out earlier.

The practical implication for ML systems is that fleet-wide failure rates can depend on age distribution. Under the bathtub model, a new cohort passes through elevated early-life risk, a relatively stable period, and eventual wear-out. Screening, operating conditions, and replacement policy determine how much of that pattern appears in field data. Mixed-age fleets combine cohorts at different stages, so their aggregate hazard need not follow the curve of a single cohort.

These three hazard regimes in figure 7.6 translate directly into admission control and maintenance policies. Before assigning new or repaired nodes to a long training run, operators use matrix-multiply, memory, and communication stress tests to expose marginal components. During service, temperature trends, error counts, and performance degradation help identify devices for scheduled replacement. Screening moves failures out of expensive production runs. An infant-mortality failure caught during a few hours of stress testing costs only idle time, while the same failure weeks into a pretraining run forces a checkpoint rollback. For that reason operators often route fresh or repaired nodes to short jobs first. Monitoring then checks whether the constant-rate model remains appropriate as the fleet ages.

7.4.3 Fleet-scale detection and mitigation

Burn-in screening and lifecycle replacement keep the most failure-prone hardware out of long training runs, but even healthy mid-life devices still suffer random faults that compound across a fleet just as the interruption rates of section 7.2 do. Field studies of DRAM in large data centers found memory errors frequent enough that error-correcting memory, background scrubbing (periodically

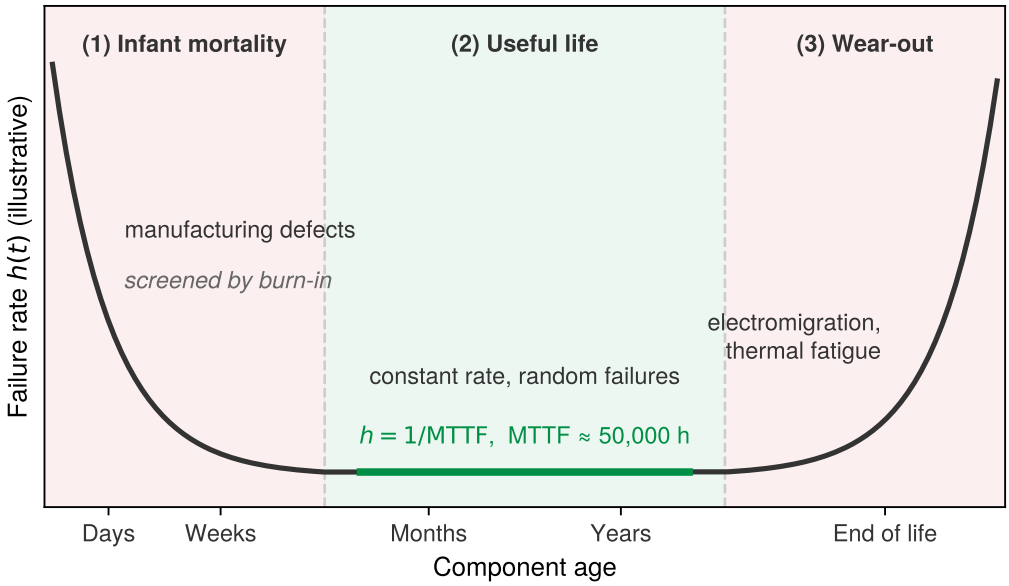

Figure 7.6: The Bathtub Curve: Hardware failure rates $h(t)$ vary over time. (1) Infant mortality: high failure rate initially due to manufacturing defects. (2) Useful life: constant, low failure rate where random failures dominate. (3) Wear-out: increasing failure rate as components age. Burn-in testing aims to filter out infant mortality failures before deployment.

rereading memory to correct errors before they accumulate), and error-rate monitoring are standard fleet practice rather than optional hardening (Schroeder et al. 2009; Sridharan et al. 2015). Protection consumes some capacity and bandwidth for check bits and scrubbing, but that cost is small compared with a corrupted multi-week run.

Detection works in layers with different coverage and cost. At the hardware layer, error-correcting codes[9] (Hamming 1950) correct and detect the memory-error patterns covered by the implemented code, and link-level CRCs[10] check corruption in transit. These protections operate within their enabled storage and transport boundaries; they do not validate an incorrect arithmetic result merely because its bytes arrive intact. Before deployment, built-in self-test and the burn-in screening of the hardware lifecycle expose manufacturing defects and permanent faults (Bushnell and Agrawal 2002). During service, operators periodically pull nodes from production and run known-answer stress tests on arithmetic, memory, and interconnect, because defective processors often fail only on particular instruction mixes, data patterns, or operating points (Hochschild et al. 2021; Dixit et al. 2021).

The software layer draws its cheapest evidence from metrics the training loop already emits. Statistical checks on loss, gradient norms, and activations flag silent drift (Chandola et al. 2009), missed heartbeats identify suspected failures (Kawazoe Aguilera et al. 1997), and watchdog timers catch tasks that stall without crashing (Pont and Ong 2002). A node that stays alive but runs slowly or erratically, the gray failure of the failure taxonomy, can degrade progress while still meeting a liveness timeout, so section 7.9.5 treats it through straggler detection. Redundant computation adds stronger evidence on selected paths, either periodically or in response to an anomaly, by replaying a step on another device or comparing independently computed results. Its frequency depends on the residual corruption risk, the value of the protected state, and the cost of detecting an error late.

[9] **Hamming codes (1950)**: Richard Hamming developed error-correcting codes at Bell Labs after relay computers repeatedly detected errors in his unattended weekend jobs and abandoned them rather than repairing them. A Hamming code uses parity checks to locate a single-bit error with $\mathcal{O}(\log n)$ check-bit overhead; adding an overall parity bit yields the extended single-error-correcting, double-error-detecting form. Memory ECC implementations vary in code and coverage, so the system must account for both corrected errors and residual faults outside that coverage.

Checkpoint 7.2: Classifying a hardware fault

Verify your understanding of the transient, permanent, and intermittent fault categories and the fault, error, failure progression:

- ☐ **Transient fault**: classify a GPU that produces a wrong result once after a cosmic-ray strike and then computes correctly for the rest of the run, and justify why detection and rollback to a known-good checkpoint, not node replacement, match that fault.
- ☐ **Intermittent vs. permanent**: distinguish an intermittent fault from a permanent one when both can produce a stuck-at value, including the diagnostic signal that separates them.
- ☐ **Fault-error-failure chain**: trace a single bit flip through the fault, error, failure chain, identifying where error correction codes and checkpoint-restart each intervene.
- ☐ **Silent corruption**: explain why silent corruption sits apart from the other fault manifestations even when its underlying fault is an ordinary transient bit flip.

7.5 Check-and-Verify Defenses Against Silent Data Corruption

Of the hardware faults just classified, the costliest are the ones whose manifestation is silent. A corrupted value that raises no error passes every liveness check the fleet runs, and training keeps building on it. Defending against silent data corruption therefore requires checks that establish a correct result, not merely a live process, and choosing those checks starts with where the corruption comes from and how it surfaces in training.

7.5.1 The physics of silent corruption

An accelerator executes deterministic logic gates, yet the silicon carrying those gates operates under narrow physical margins where both random particle strikes and load-dependent defects can flip bits without halting the processor. Single-event upsets (SEUs), the particle strikes introduced with transient failures, create fleet exposure that grows with device count and operating time even when the per-device rate is low. Manufacturing variances can instead leave "marginal" chips that pass initial QA but fail under particular inputs or operating conditions. Voltage and temperature margins matter, including supply disturbances during rapid current changes (di/dt). Because these margin defects produce repeatable errors under specific workloads while particle strikes arrive at random, fleet reliability requires a probabilistic exposure model even when individual hardware faults are deterministic.

Facebook (now Meta) documented a pervasive SDC issue where a hardware fault caused a valid file to be reported as "size zero" during decompression (Dixit et al. 2021); the system "worked" without crashing, but data was silently deleted. In ML, the same failure manifests as valid-looking but numerically corrupted gradients. The pipeline trace in figure 7.7 follows a single ALU or memory bit flip as it corrupts an individual tensor entry, taints gradient reduction during the backward pass, commits into weights, and finally surfaces as degraded quality at evaluation. The comparative timelines beneath contrast this silent latency with a standard detected error. A hardware trap halts within milliseconds to restart cleanly from the last checkpoint, whereas silent corruption commits contaminated state into checkpoints over hours or days before an evaluation alert forces a costly rollback.

Telemetry from production training fleets shows this corruption is routine rather than anomalous. At modern training scales, hardware bit flips occur regularly, and erroneous computations from a single malfunctioning accelerator can propagate into the global model state, as Jeff Dean documented in his MLSys 2024 keynote on silent data corruption (Dean 2024) (figure 7.8).

In the Google training telemetry presented by Dean (2024), figure 7.9 identifies an abrupt gradient-norm anomaly as an SDC event. The plot shows the symptom without identifying the damaged bit or circuit. Corrupting a floating-point exponent is one mechanism that can greatly change a tensor value and subsequent updates. Gradient clipping can limit update magnitude, but it does not reconstruct the correct gradient; an anomaly detector must trigger validation or recovery before suspect state is accepted.

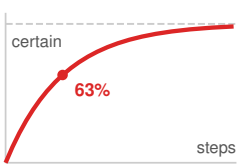

Cumulative SDC probability nears certainty over a long run.

7.5.2 Verification at fleet scale

Even a small per-GPU SDC rate accumulates across repeated collective operations, and the accumulation grows with fleet size. Any single collective is unlikely to contain an SDC event. Under

[10] **CRC (cyclic redundancy check)**: Polynomial checksum family introduced by Peterson and Brown (1961). CRC coverage depends on the polynomial, frame length, and error model; it is best understood as a low-cost detection layer rather than a universal guarantee. In distributed ML training, checksums or hashes can validate gradient payloads exchanged during collectives; without a verification layer, a corrupted gradient packet can silently poison the parameter update for every worker in the collective.

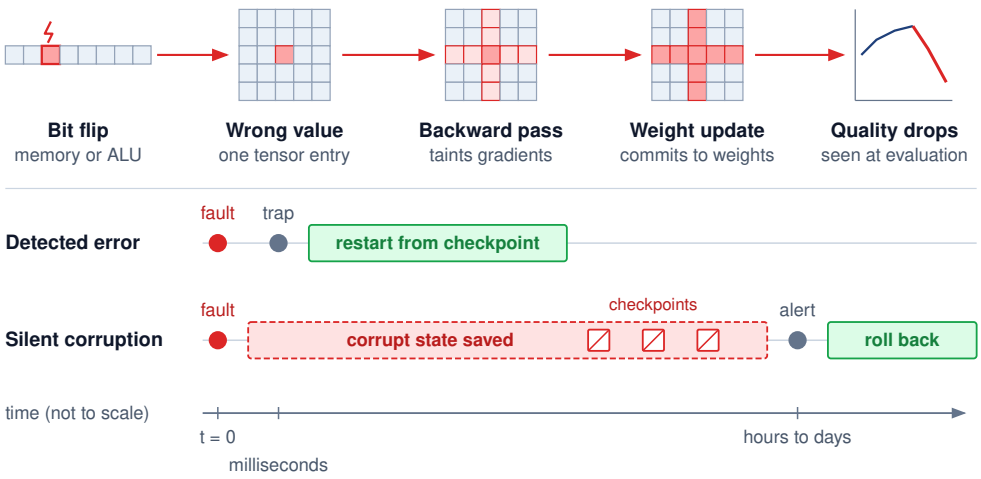

Figure 7.7: Silent Data Corruption Propagation: A hardware bit flip moves left to right from one wrong tensor value through the backward pass and the weight update, surfacing only at the final stage as degraded model quality. The two timelines below contrast a detected error, which traps within milliseconds and restarts from the last checkpoint, with silent corruption, which is saved into checkpoints for hours or days before an evaluation alert forces a rollback; the gap between the two is the cost of silence.

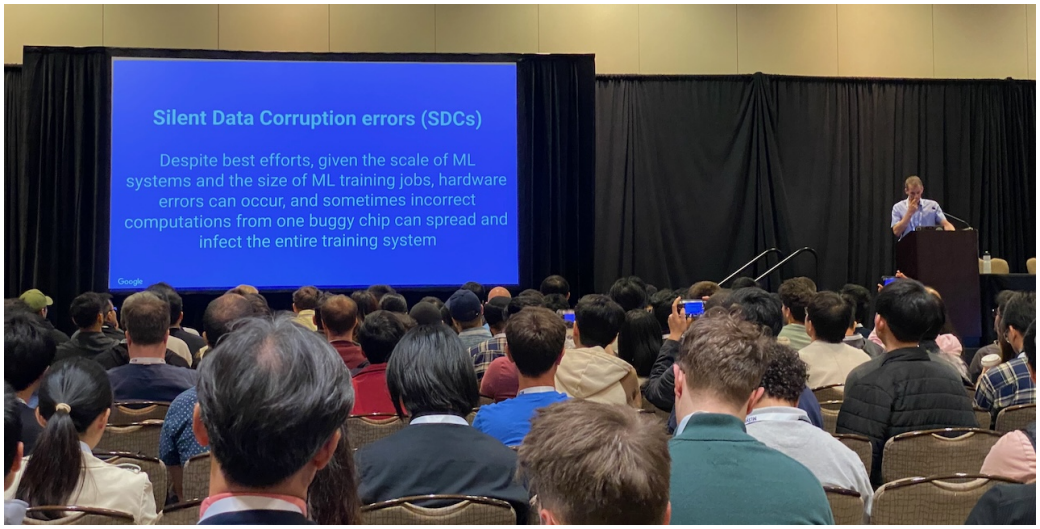

Figure 7.8: Silent Data Corruption at Scale: Jeff Dean presents on silent data corruption (SDC) at MLSys 2024, arguing that at the scale of large ML training jobs hardware errors occur routinely, and that incorrect computations from one buggy chip can propagate and infect an entire training run. Source: Jeff Dean, MLSys 2024 keynote (Google).

the illustrative per-GPU rate stated in the notebook below, the per-step probability remains small while the cumulative probability grows over a long training run, reaching $1 - e^{-1} \approx 63$ percent once the expected number of events reaches one. On the chapter's 10,000-GPU cluster, that rate gives one expected event roughly every 100 hours, or about 22 over the three-month run of the chapter opening. The notebook below scales the same estimate to a fleet ten times larger to show where the trend leads, since the expected interval shrinks in proportion to fleet size.

Standard AllReduce reduces the supplied contributions; it does not establish that an *alive* worker computed a *correct* gradient. The fleet therefore needs validation appropriate to its fault model in addition to liveness checks. The robust aggregation rules of section 7.3.3, Krum and trimmed

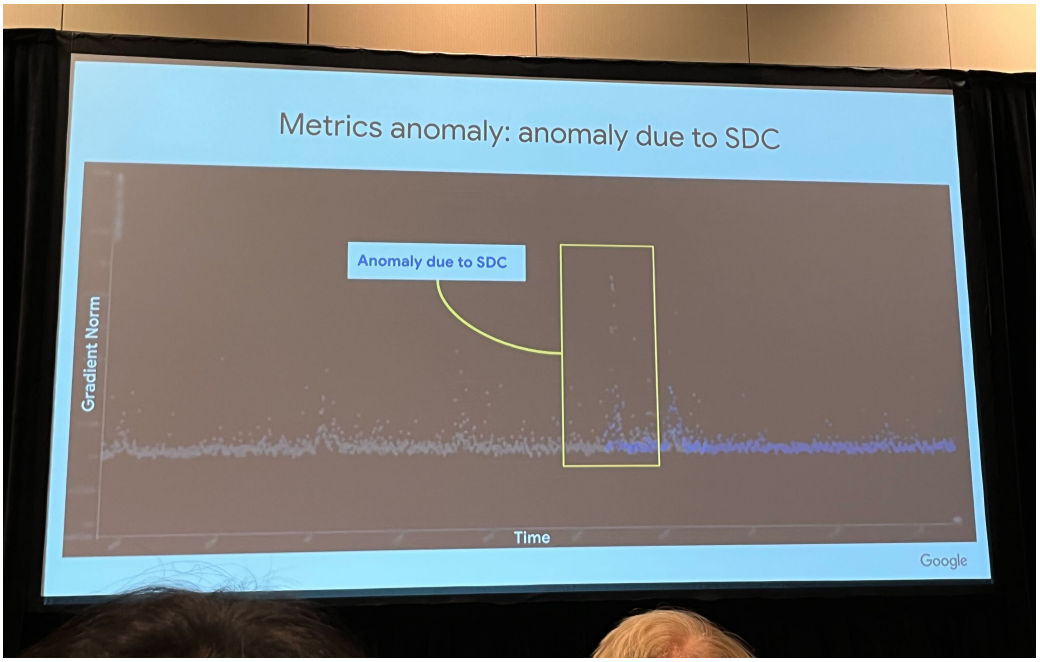

Figure 7.9: Gradient Norm Deviation: Transient hardware faults, such as silent data corruption (SDC), can disrupt optimization by causing abrupt changes in gradient norms. Google's production examples show SDC anomalies surfacing as visible gradient-norm spikes. Source: (Dean 2024).

mean, tolerate a bounded number of wrong gradients at extra communication and computation cost. Check-and-verify methods instead look for the wrong values, testing data-path invariants, subject to the boundary drawn for transient failures in section 7.3.1. Digests test byte integrity and agreement within their error-detection limits; they cannot establish that the original computation was correct. Independent computation, an algorithmic invariant, or replay tests that additional claim under a declared fault model. A quick estimate shows how much exposure those checks must cover.

Napkin Math 7.1: SDC risk accumulation

Problem: Under an illustrative generic SDC-rate assumption, calculate the probability that at least one GPU in a 100,000-GPU fleet experiences an SDC event during a single 2 s training step.

1. **Fleet size**: 100,000 accelerators.
2. **Individual risk**: 10^{-6} per GPU-hour (an illustrative generic SDC-rate assumption).
3. **The exposure**: In a 2-second window, the fleet has $100{,}000 \times (2/3600) \approx 56$ "GPU-hours" of exposure.
4. **The probability**: $\Pr(\text{at least one SDC}) \approx 0.0056$ percent.

Systems insight: Under this illustrative constant-rate assumption, a 100k-GPU fleet expects one SDC event every 18,000 steps (roughly every 10 hours), even though the per-step probability is only 0.0056 percent. A CRC or hash can protect each shard or reduced buffer against corruption in transit or storage, while a redundant reduction can test the arithmetic path independently. Each check moves robustness from a *restart* problem toward a *verification* problem only for the failure mechanism it covers.

Full replication protects a broader computation path at a correspondingly high resource cost. Dual modular redundancy runs every computation twice and compares the results, and triple modular redundancy runs it three times and takes a majority vote. These schemes pay 100 and 200 percent

hardware overhead, respectively, to detect or mask a fault within a single control deadline (Sheaffer et al. 2007; Arifeen et al. 2020). These guarantees depend on the replicas avoiding a shared fault. The price is justified in safety-critical embedded control, such as a vehicle's onboard computers, where one wrong output can endanger people and the replicated hardware is a small computer. For fleet training, duplicating every step of a run on 10,000 GPUs doubles the compute devoted to those steps, before comparison and coordination costs. When corruption is rare and rollback can recover valid state, selective redundancy can protect useful progress at lower cost. Screening removes suspect devices, statistical checks run continuously, standby pools absorb drained nodes, and periodic or triggered recomputation checks selected execution paths.

Google's reactive hot-spare protocol shows how those pieces combine in three stages (figure 7.10). When runtime monitors flag an anomalous gradient and replay, which reruns the flagged computation on other hardware and compares the results, implicates a suspect accelerator (G2), the scheduler isolates that device and promotes a standby hot spare (S) into the active ring. It then provisions a replacement spare (S') while the suspect unit undergoes hardware diagnostics (Dean 2024). The fleet pays in advance for one idle device per pool rather than for a duplicate of every computation, and it gives up coverage in exchange. The spare replaces only a device that monitors and replay have flagged, so corruption those checks miss stays in the training state. The protocol shifts the burden of correctness from unverified hardware to system-level auditing. Even so, these hardware checks reduce only the risk of an unreliable physical substrate; they do not protect against faults that the software stack introduces itself, bugs that look like valid computation while corrupting the ML pipeline.

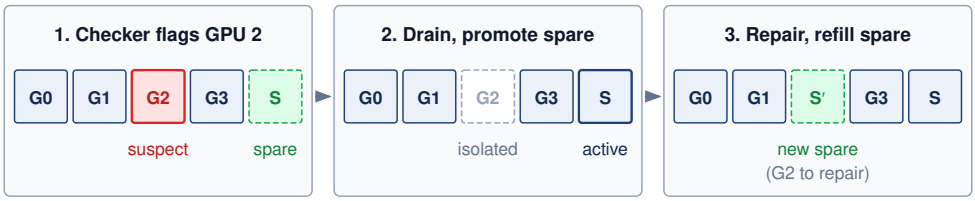

Figure 7.10: Hot-Spare Failover: An $N+1$ pool of four active GPUs and one standby spare. When loss or gradient sanity checks flag a suspect GPU and replay implicates it, the job drains that device and resumes on the spare from the last known-good state; the device goes to repair and a healthy unit refills the spare slot. Google's data-center examples use this reactive pattern to keep training moving despite hardware faults, in contrast to running every computation on two or three redundant units and comparing results (Dean 2024). The standing cost is the idle spare.

Checkpoint 7.3: Detecting silent corruption

Verify your understanding of Byzantine failures and SDC:

- ☐ **SDC danger**: explain why silent data corruption is more dangerous for ML training than a simple node crash.
- ☐ **Hot spare vs. restart**: compare a hot spare redundancy strategy with traditional checkpoint-restart in terms of recovery latency.
- ☐ **Sanity checker**: identify the metric used as a "sanity checker" to identify potential hardware faults: GPU temperature or gradient norm.
- ☐ **Byzantine scope**: assess the True or False claim that Byzantine failures only occur when a malicious actor intentionally sabotages the training cluster.

7.6 Software Faults

A team spends three months testing a large language model (LLM) against complex prompt-level failure cases, only to realize that its preprocessing script kept only the last 512 tokens of each input, silently discarding the system prompt on long requests. The example exposes an ordinary source of

ML failure that algorithmic testing can miss. A logic error in a data loader or prompt assembler can preserve valid tensor shapes while changing the computation's meaning, allowing the pipeline to continue without a crash.

Software faults require a different recovery lens from hardware failures. Hardware faults enter from silicon, electrons, and physical wear; software faults enter through design choices, dependency versions, and pipeline glue. The fault-tolerance problem shifts from replacing a failed component to identifying which valid-looking transformation corrupted data, gradients, or predictions. Every boundary between framework, library, runtime, and distributed system is a place where such a transformation can hide.

7.6.1 Software fault properties and propagation

A single tensor reference leaked on every training iteration lets a correctly sized job drift toward the device memory limit over hundreds of steps and then crash with an out-of-memory error far from the code that leaked it. Across ML frameworks, software faults of this kind span logic errors, resource leaks, concurrency bugs, and integration mismatches that propagate across system boundaries. Race conditions in asynchronous updates can lose or overwrite parameter updates, and deadlocks[11] leave workers waiting on messages that never arrive. Because many of these faults surface only under high load, on particular hardware, or with a specific pairing of a high-level framework and a low-level numerical library such as cuDNN, root-cause analysis must treat the model code, the runtime, and its dependencies as one versioned artifact.

11 | **Pipeline deadlock**: In pipelined training, a deadlock can form when stage 1 waits for a free buffer to send activations to stage 2 while stage 2 waits for stage 1 to accept its backward gradients. Neither stage can proceed, and because every other stage depends on them, the whole job stalls with no error raised, which a collective timeout eventually converts into a detectable failure.

7.6.2 Software fault detection and prevention

The prompt-truncation bug from the opening example and the memory or synchronization bugs above should never survive unchallenged into a multi-week run. Because many software faults leave the pipeline running while corrupting inputs or gradients, each verification gate must establish which invariant holds and which recovery action is safe when it fails. A unit test catches the tokenizer boundary before launch, an integration replay verifies that a recorded batch reproduces the expected shapes and state on known-good hardware, and a fault-injection test verifies that a corrupted buffer or lost worker triggers rollback before tainting committed state.

Several of these gates also test the recovery machinery itself before an incident relies on it. A restore drill verifies a checkpoint manifest, the record that names every saved shard and the training step they share (section 7.8.5.2). A membership fault-injection test kills or partitions a worker and checks that, once replaced, the stale process can no longer write state the job will trust. The epoch, lease, and fencing rules of section 7.9.2 enforce that guarantee by stamping every write with the current membership counter, so storage rejects any write that carries an older one. Every recovery drill is judged against two bounds formalized in section 7.8: how much completed work a failure may lose and how long recovery may take before useful execution resumes. Across all six gates, each test earns its place through the operational decision it authorizes (table 7.4).

Table 7.4: Reliability Evidence and Recovery Decisions: Tests earn operational value when each one establishes an invariant and authorizes a specific rejection, rollback, isolation, or restore decision.

Reliability evidence	Invariant exercised	Decision enabled
Boundary unit test	Tokenizer limits, tensor shapes, masks, and step metadata remain valid	Reject the change before distributed execution
Deterministic replay	A recorded batch produces equivalent state on known-good hardware	Separate input or software defects from suspect hardware
Checkpoint restore drill	A committed manifest restores every shard and resumes at its named step	Accept the generation as satisfying the recovery contract
Membership fault injection	A killed or partitioned worker loses ownership and cannot write after replacement	Validate ownership handoff and communicator rebuild
Corruption injection	A changed shard, reduced buffer, or checkpoint fails the intended verifier	Block the optimizer step or roll back before propagation
Dependency rollback test	A versioned runtime and artifact set can return to the last known-good combination	Authorize automated rollback instead of manual repair

These safeguards matter only when each gate targets a concrete corruption path. A test that checks only whether the tokenizer runs is too weak; it must assert that system prompts, labels, masks, and sequence-length metadata survive preprocessing with the intended semantics. Integration replay then follows one recorded batch through loading, sharding, collation, and model input construction. Fault injection tests whether the detector fires before the protected commit and whether recovery stays within the named bounds on lost work and recovery time for the injected fault. Restore drills, stale-writer rejection, and corrupted-shard rejection therefore provide evidence that generic build success cannot supply, and section 7.14 exercises those claims directly once the recovery paths themselves have been built. Before building those paths, however, a fleet must decide how much protection each workload is worth.

7.7 The Cost of Failure

Losing an uncheckpointed day on a 10,000-GPU foundation-model run burns hundreds of thousands of dollars in accelerator time, whereas losing an hour on a small vision job costs a few hundred dollars and a failed update in an online recommender costs live revenue rather than batch rework. Hardware, silent corruption, and software together set how often a fleet fails, but the work a failure destroys determines what each incident costs. A defense is worth buying only where it removes more of that expected loss than it adds in standing overhead.

7.7.1 Failure cost by workload

Three workload properties set the price of a training interruption: the burn rate of the cluster running the job, the footprint of state that must be saved and reloaded, and how sensitive model quality is to lost or reordered updates. Those three factors separate multi-terabyte foundation models and continuous recommenders from smaller vision, speech, and scientific workloads (table 7.5).

Table 7.5: Failure Impact by Model Type: The cost of training failures varies dramatically by model type, driven by training duration, checkpoint overhead, and the value of accumulated training state. The Archetype A cost is one lost day on the chapter's 10,000-GPU cluster; the vision, speech, and scientific rows give illustrative ranges rather than measured values. These differences demand model-specific fault tolerance strategies.

Model type	Typical training duration	Checkpoint size	State sensitivity	Failure cost
Archetype A (GPT-4/Llama-3)	Weeks to months	2.45 TB (175B)	High (position in data order)	$0.5M compute per 24 hr loss
Vision (ViT-Large)	1–3 days	1.2 GB	Medium (augmentation state)	$10–50K per day loss
Archetype B (DLRM at Scale)	Continuous	10 TB+ (embeddings)	Very High (embedding freshness)	Revenue impact per hour
Speech (Whisper-scale)	3–7 days	5–10 GB	Medium	$50–200K per day loss
Scientific (AlphaFold)	Days to weeks	10–50 GB	High (exploration state)	Research delay

Large language model training illustrates how a costly allocation magnifies each hour of lost progress. The chapter's 10,000-GPU cluster at approximately $2/GPU-hour incurs $0.5M in compute costs per day. A job checkpointed only once a day can lose up to a full day of progress to one failure, which costs $0.5M in wasted compute plus schedule delay. At the checkpoint interval derived in section 7.8.1, the same failure costs minutes of rework instead of a day. The checkpoint overhead is correspondingly large. At the 14-byte-per-parameter resumable format that the checkpointing analysis uses (BF16 weights, FP32 master weights, and two FP32 Adaptive Moment Estimation (Adam) moments), the 175B-parameter running example writes 2.45 TB at every save.

7.7.2 Investing in protection

Every GPU cycle, network byte, and storage drive dedicated to checkpointing, standby nodes, or verification is a resource that is not training the model, yet skimping on protection leaves a multi-million-dollar run exposed to day-long rollbacks. Rational investment in fault tolerance therefore weighs what a failure destroys against what the defense costs to run.

When a failure strikes, its bill extends well beyond the GPU-hours burned on steps that must be replayed. Repeating lost work delays the delivery of the trained model, ties up accelerators that could have run other jobs, and pulls engineers away from research to diagnose faults and nurse manual restarts. On the other side of the cost account, prevention carries a continuous overhead even when the cluster is healthy. Checkpoint storage scales with model size and save frequency, and checkpoint writes compete for memory and network bandwidth. Hot spares and automated recovery controllers consume standing infrastructure, and fault-tolerant coordination adds software complexity that must be maintained and debugged.

Optimal investment balances these two curves against cluster scale. For a small job on a few GPUs where hardware failures are weeks apart, infrequent checkpoints and manual restarts are cheaper than standing recovery infrastructure. On thousands of GPUs where failures arrive several times a day, frequent checkpoints, automated failover, and elastic resizing pay back their overhead many times over. The checkpoint interval derived in section 7.8.1 applies this exact marginal balance to a single decision variable, choosing the cadence where the cost of one more save equals the expected rework it prevents.

Equation 7.5 presents a simplified economic model for expected cost per training run:

$$C_{\text{total}} = C_{\text{compute}} + \mathbb{E}[N_{\text{failures}}] \times C_{\text{per-failure}} + C_{\text{ft}} \tag{7.5}$$

where C_{compute} is the base compute cost, $\mathbb{E}[N_{\text{failures}}]$ is the expected number of failures during training, $C_{\text{per-failure}}$ is the cost per failure, and C_{ft} is the cost of fault tolerance mechanisms. The cost per failure includes wasted compute plus overhead.

Equation 7.6 formalizes when fault tolerance investment is justified:

$$\frac{\partial C_{\text{ft}}}{\partial x} < -\frac{\partial(\mathbb{E}[N_{\text{failures}}] \times C_{\text{per-failure}})}{\partial x} \tag{7.6}$$

where x is the amount of protection bought, a design variable such as checkpoint frequency or the number of hot spares, and the base compute cost is held fixed. The negative derivative on the right is the marginal reduction in expected failure cost. Additional protection reduces total expected cost while this reduction exceeds its marginal cost; an interior optimum equates the two when the cost functions have a suitable minimum.

One investment worked through the criterion shows what it decides. Let x be a hot-spare node held idle beside the chapter's 10,000-GPU cluster, so that a failed node's replacement no longer waits for the scheduler to find capacity. Its marginal cost is 8 GPUs idling all day, about $384 per day. Its marginal benefit is the recovery time it removes from each failure while every other GPU waits. At the cascade's 6.2 failures/day, each minute removed from every recovery is worth about $2,100 per day, so the spare pays for itself once it shortens each recovery by about 11 seconds, provided it is refilled before the next failure arrives some 3.9 hours later. The same node beside an eight-GPU job, whose GPU-only MTBF in table 7.1 is 6,250 hours (260 days), would sit idle far longer than it could ever save. Scale reverses the verdict on the same hardware. That reversal reflects the first of three rules of failure at scale.

> **Systems Perspective 7.2: Three rules of failure at scale**
>
> 1. *At scale, failures are continuous, not exceptional.* A 10,000-GPU cluster experiences a failure every 3.9 hours on average. Systems must be designed expecting failure as normal operation.
> 2. *Checkpoint intervals have an optimum.* The Young-Daly interval (developed in section 7.8.1) sets checkpoint frequency from write time and system MTBF under its stated assumptions, and it shrinks only with the square root of a falling MTBF.

3. *Training and serving have fundamentally different fault tolerance requirements.* A long training job can budget minutes for recovery while preserving accumulated progress; an interactive serving request may have only milliseconds left before its deadline. Those budgets motivate checkpoint recovery for training and ready alternate capacity for serving.

These three rules organize the rest of the chapter. Rules 1 and 2 govern long-running training jobs that must protect accumulated progress, rule 3 governs low-latency serving fleets that must protect live requests, and observability and fault injection verify that both recovery paths work under fire. Training resilience begins with the durable state that every restart or resize falls back on.

7.8 Checkpointing and Progress Preservation

When a top-of-rack switch dies two weeks into the training run of a 175-billion-parameter model, the cluster loses 64 GPUs instantly. Because a large cluster experiences such interruptions routinely (section 7.1), it cannot restart multi-week jobs from scratch; training fault tolerance must preserve state, detect and classify the failure, and resume or resize the worker group. Checkpointing supplies the first of those three obligations, and every detection, restart, and elastic-resizing path later in the chapter depends on the durable snapshot it produces. The first design question is what that snapshot must contain to make a restart valid.

Definition 7.1: Checkpointing

Checkpointing is the periodic serialization of the complete training state (parameters, optimizer state, data loader position, and random state) to persistent storage at a recorded training step.

1. **Significance**: It minimizes lost work after a system failure by trading checkpoint I/O against rework. Writing too often wastes accelerator time on storage traffic; writing too rarely risks replaying many completed steps after a failure. The Young-Daly interval developed in this section chooses the cadence that balances write cost against expected lost work.
2. **Distinction**: A resumable checkpoint captures a consistent execution state, including optimizer state, random-number generator state, data position, and learning rate schedules. It can use full or incremental storage; bitwise replay additionally requires compatible deterministic execution.
3. **Common pitfall**: A frequent misconception is that checkpointing is "just writing to disk." For large models it is a storage-system stress test, because the simultaneous write from thousands of GPUs can trigger a checkpoint storm (definition 4.5) that saturates the shared network and storage bandwidth (BW).

A checkpoint is useful only if it resumes the recorded step consistently. In distributed training, every shard must belong to the same recoverable state; a collection of files from different steps does not preserve progress. For the interval model, assume that writing a checkpoint completes durably before it becomes available for recovery and that a failure discards the work performed since the last completed checkpoint.

Two recovery objectives turn the requirements of a valid restart into a contract. The recovery point objective (RPO) bounds acceptable lost progress, so checkpoint age and replication lag must fit within it. The recovery time objective (RTO) bounds the delay until useful work resumes, which covers detection, resource replacement, restore, and warmup. A 20-minute checkpoint interval cannot promise an RPO below 20 minutes without a more frequent durable log or replica, and asynchronous flushing can leave the last durable generation older still. A five-minute RTO is impossible if checkpoint reload alone takes eight minutes. These objectives apply to a stated fault model. Corruption discovered after several saves can require rollback beyond the newest durable checkpoint. The Young-Daly interval derived next minimizes expected waste; RPO and RTO impose service constraints that the expected-value optimum must still satisfy (Beyer et al. 2016).

7.8.1 The Young-Daly checkpoint interval

When failures are expected during a job, the key engineering decision is how often to write that state, because every save costs I/O time and every unsaved step is work a failure can erase. The Young-Daly formula[12] (principle 13) gives a first-order estimate of that balance through a square-root rule:

$$T^*_{\text{ckpt}} = \sqrt{2 \cdot T_{\text{write}} \cdot \text{MTBF}_{\text{system}}} \tag{7.7}$$

T_{write} in distributed ML training is not the time to flush a small process state; it includes transferring model weights and optimizer state to durable storage. The reference format here stores BF16 weights, separate FP32 master weights, and two FP32 moment tensors per parameter for Adam. A single resumable checkpoint for a 70B-parameter model therefore approaches a terabyte at 14 bytes per parameter; for the later 175B-parameter running example, it exceeds two terabytes. Of those 14 bytes per parameter, the optimizer state[13] dominates the payload, exactly as in the storage sizing of the same resumable checkpoint in section 4.6, so the optimizer state rather than the model weights sets most of T_{write}. Reducing T_{write} requires a smaller payload or higher sustained write bandwidth, and T_{write} is the central constraint that determines how tightly the checkpoint interval can be set before I/O overhead displaces productive training compute.

To see the balance, consider an interval between saves. Doubling that interval approximately halves the checkpoint writes per unit of training, but a failure then loses twice as much work on average. Under the constant-rate, rare-failure approximation, failure occurs approximately uniformly within an interval, so the average lost work is half the interval. When write time is small relative to the interval, the wasted fraction from saving and replaying work is approximately the sum below. Detection, allocation, and restoration time require a separate recovery budget.

$$W(T_{\text{ckpt}}) \approx \frac{T_{\text{write}}}{T_{\text{ckpt}}} + \frac{T_{\text{ckpt}}}{2 \cdot \text{MTBF}_{\text{system}}} \tag{7.8}$$

Setting the first derivative to zero ($\frac{dW}{dT_{\text{ckpt}}} = -\frac{T_{\text{write}}}{T^2_{\text{ckpt}}} + \frac{1}{2 \cdot \text{MTBF}_{\text{system}}} = 0$) directly yields the square-root minimum of equation 7.7. As clusters grow, system MTBF drops, so the optimal interval shrinks, demanding higher-bandwidth storage to keep T_{write} small before checkpoint overhead consumes the cluster's compute capacity. Because the relationship is a square root, halving MTBF tightens the interval by only $\sqrt{2} \approx 1.4\times$, so doubling cluster size does not demand doubling checkpoint I/O bandwidth when checkpoint size and the other assumptions stay fixed. Substituting the optimum back into equation 7.8 gives a minimum waste of $\sqrt{2T_{\text{write}}/\text{MTBF}_{\text{system}}}$. Because $\text{MTBF}_{\text{system}}$ falls as $1/N$ under independent failures, that optimized overhead grows as $\sqrt{N}$, and doubling the cluster raises it by a factor of $\sqrt{2}$ even though the interval itself adapts.

7.8.2 Checkpoint interval from failure analysis

The Young-Daly interval of equation 7.7 needs two inputs, the system MTBF and the write time T_{write}, and the chapter's 10,000-GPU cluster supplies both as computed values. The worked cascade (section 7.2.2) put full-system MTBF, including PCIe and network failures (the redundant power supplies add no term), at 3.87 hours. For the write time, the 175B running example's 2.45 TB resumable checkpoint drains to shared storage at the sustained aggregate write bandwidth this chapter assumes for the cluster, 100 GB/s, which takes 24.5 s. Substituting these into the Young-Daly formula gives a candidate interval for this cluster, which must still satisfy its RPO and RTO.

> Napkin Math 7.2: The Young-Daly optimal interval
>
> **Problem**: A 10,000-GPU cluster has an MTBF of 3.87 hours. A full model checkpoint takes 24.5 s to write. What is the optimal checkpoint frequency?
>
> **Math**: Apply the Young-Daly interval of equation 7.7 (principle 13).
>
> 1. **Convert to common units**: $\text{MTBF}_{\text{system}} = 3.87$ hours $\times$ 3600 s/hour = 13932 s.

[12] **Young-Daly formula**: Young (1974) derived the first-order optimal checkpoint interval, and Daly (2006) independently refined it with tighter second-order bounds.

[13] **Adam optimizer state**: Adam's per-parameter first-moment (m) and second-moment (v) estimates are two FP32 values, 8 of the 14 bytes per parameter in the resumable checkpoint; with the FP32 master weights, optimizer-side state accounts for all but the 2 bytes of BF16 weights. For a 175B parameter model, the two moments alone reach 1.4 TB, which typically dominates checkpoint size and recovery time, making optimizer state the primary bottleneck in checkpoint I/O design.

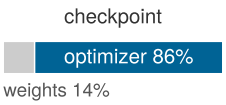

Optimizer state, not weights, dominates checkpoint size.

↩ The per-tier storage bandwidths that bound how quickly a checkpoint can be written are established in section 4.3.

2. **Substitute and solve**: With $T_{\text{write}} = 24.5$ s and $\text{MTBF}_{\text{system}} = 13932$ s, the formula gives $T^*_{\text{ckpt}} \approx 826.2$ s, or about 13.8 min.

Systems insight: At this interval, the checkpoint overhead (time spent saving plus time spent recomputing) is minimized to approximately 5.9 percent. Checkpointing every hour instead would lengthen the expected rework after each failure, wasting about 13.6 percent of the cluster's capacity.

The result connects measured or estimated interruption rates to checkpoint policy. At this interval, saving and expected replay together consume approximately 5.9 percent of training time in the first-order model, half from writes and half from replay, because at the square-root optimum the two terms of equation 7.8 are equal. Plotting those two terms for this cluster shows how forgiving the optimum is (figure 7.11). Write overhead falls hyperbolically and expected rework climbs linearly, so their sum has a shallow floor around 14 min. An interval slightly off the optimum costs little, while one several times too short or too long costs far more, which is why a rough MTBF estimate still yields a usable schedule. How far the estimate itself can be trusted depends on the simplifying assumptions behind it.

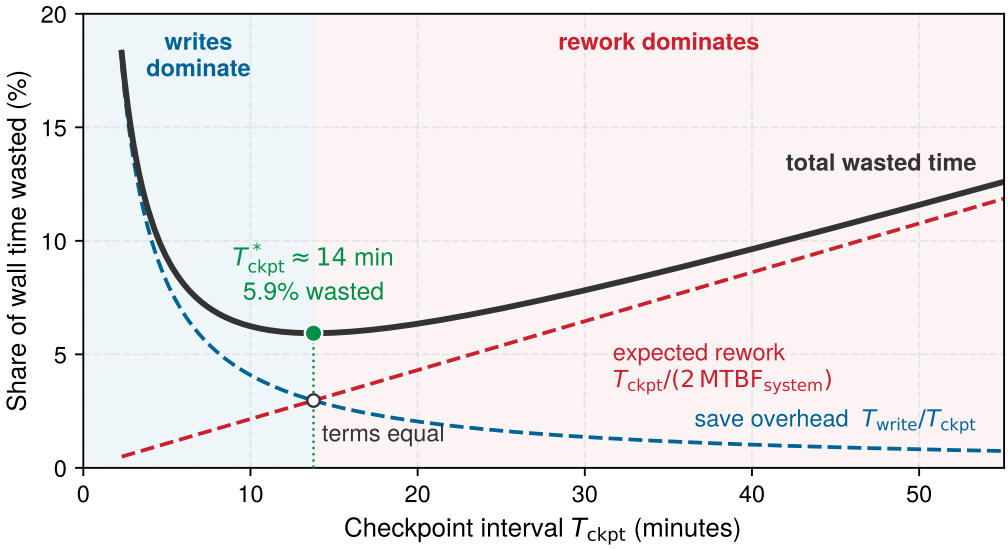

Figure 7.11: The Checkpoint Trade-Off: Checkpoint save overhead (dashed blue) falls as the interval lengthens, while expected rework after a failure (dashed red) rises; their sum (solid) is minimized at the Young-Daly interval. The curves use this chapter's 175B running example on its 10,000-GPU cluster, with a 24.5 s checkpoint write and a 3.9 h full-system MTBF, which put the optimum near 14 min. The minimum is shallow, so an interval somewhat off the optimum costs little, while a far shorter or far longer one costs a great deal.

Systems Perspective 7.3: Young-Daly formula assumptions

The first-order Young-Daly formula provides valuable intuition but rests on assumptions that may not hold when write times are non-negligible relative to MTBF; the renewal model derived in section E.5 shows how finite checkpoint time changes the optimum. The assumptions are:

- Exponentially distributed failures: Assumes a constant failure rate. Real hardware follows the bathtub curve, with higher rates during burn-in and wear-out phases.

- Deterministic checkpoint time: Assumes T_{write} is constant. In practice, checkpoint time varies from save to save with storage contention, network congestion, and memory pressure.
- Recovery cost is independent of the chosen interval: A fixed per-failure recovery cost does not change the first-order minimizing interval. Recovery can take longer than checkpoint writing because it also includes job scheduling delays, topology reconstruction, and warmup; it still belongs in the total-waste budget.
- Aggregate interruption process: Independent constant-rate failure sources can contribute to one job-interruption rate. Shared power, cooling, or switch failures require a model of incident frequency and recovery cost; adding per-device rates alone misses those common causes.
- Infinite timeline: Optimizes for long training runs. Short runs, where total time is comparable to MTBF, require different analysis.

When these assumptions are violated, the optimal interval may shift significantly. Restart time adds a per-failure term to total expected waste rather than a cost paid at every checkpoint; it does not replace T_{write} inside equation 7.7 unless a recovery-aware Daly-style model is explicitly derived and calibrated.

Along an execution timeline, a training run alternates productive compute intervals (T_{ckpt}) with synchronous save barriers (T_{write}), and a failure charges two distinct penalties (figure 7.12). Uncheckpointed progress is discarded as lost rework, and the restart latency (T_{recovery}) of job scheduling, checkpoint loading, and pipeline warmup follows before useful work resumes. Restart adds scheduling and warmup to reading the saved state back, so total restart time can exceed the checkpoint write time. Recovery latency belongs in the total-waste and service-level objective (SLO) budgets as a per-failure term (perspective 7.3) rather than inside the per-checkpoint write term.

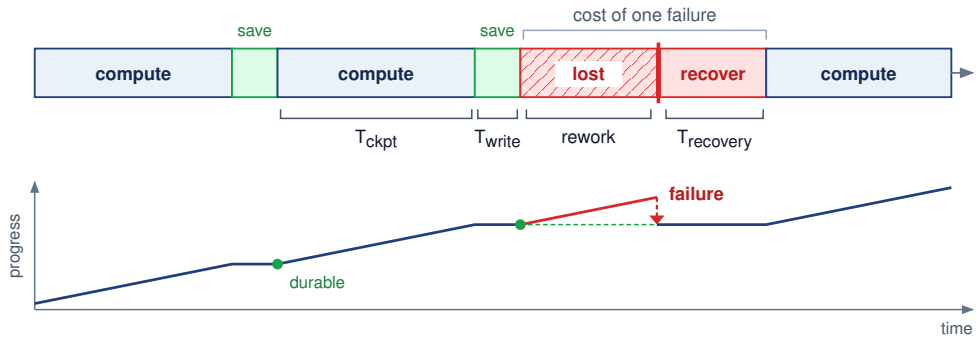

Figure 7.12: Checkpoint-Recovery Timeline: A training run alternates productive compute (blue) with durable checkpoint writes (green). When a failure strikes, all compute since the last completed checkpoint is lost (red hatching), and the job must then read the checkpoint and restart (red) before productive compute resumes. A failure therefore costs the rework since the last save plus the recovery latency; the Young-Daly interval trades the first of these against the write time paid at every checkpoint.

7.8.3 Checkpoint overhead analysis

The interval calculation chooses *when* the system preserves state; the write path decides what each save costs while training waits. Equation 7.9 generalizes the save-side term of equation 7.8 to include any pause beyond the write itself:

$$f_{\text{ckpt}} = \frac{T_{\text{write}}}{T_{\text{ckpt}}} + \frac{T_{\text{pause}}}{T_{\text{ckpt}}} \tag{7.9}$$

Here, f_{ckpt} is the dimensionless fraction of time lost to checkpointing, T_{write} is checkpoint write time, T_{ckpt} is the checkpoint interval, and T_{pause} is any training pause beyond the checkpoint write itself. This includes memory allocation, coordination, and serialization. Total first-order waste is then $W(T_{\text{ckpt}}) \approx f_{\text{ckpt}} + T_{\text{ckpt}}/(2 \cdot \text{MTBF}_{\text{system}})$, so checkpoint overhead in this section means the save side alone, with expected rework counted separately. How large f_{ckpt} becomes depends first on whether training waits for durable storage.

7.8.3.1 Synchronous vs. asynchronous checkpointing

Synchronous checkpointing resumes training only after every worker's portion reaches durable storage. Asynchronous checkpointing instead pauses the accelerators only long enough to copy the weights and optimizer tensors from GPU memory into host memory or local staging, then flushes that staged copy to remote storage in the background while training continues. The brief device-to-host copy is necessary because the next optimizer step updates the live GPU tensors in place, which would otherwise overwrite state that is still being serialized. Section 4.6.1 develops the write path behind both approaches. For recovery, blocking alone does not establish consistency, so the protocol must ensure that every shard belongs to the same recoverable training state. Recovery can use an asynchronous generation only after all required shards reach the storage tier that survives the modeled failure and the coordinator publishes completion. If a failure strikes while the background flush is still in flight, recovery must fall back to the previous committed generation.

A synchronous save pauses every accelerator in the job for the full remote write, and section 4.6 prices that exposed pause on a slow shared tier. Asynchronous capture shrinks the exposed pause to the local staging copy and moves the remote flush off the critical path, which is why production systems stage checkpoints rather than block on shared storage. The Young-Daly interval computed earlier from a blocking write is therefore the conservative case; a staged save that exposes a shorter pause shifts the optimum toward shorter intervals and lower save overhead, provided each background flush completes within the interval. Write time also varies widely across model types (table 7.6). Larger models require longer writes but also earn longer optimal intervals, so save overhead grows only with the square root of write time.

The Archetype A row in table 7.6 uses the same 14-byte-per-parameter resumable checkpoint (2.45 TB) as the rest of the chapter. Its interval calculation uses the GPU-only cluster MTBF rather than the shorter full-system MTBF of the worked cascade, so the square-root rule gives a longer interval than the Young-Daly notebook. The table keeps the GPU-only basis because its purpose is comparison across model types. Switching every row to the full-system MTBF would shorten each interval, and raise each overhead, by the same square-root factor, leaving the ranking across rows unchanged.

The ResNet-50 and BERT-Large rows count only FP32 weights, so their write times and resulting intervals are lower bounds for a resumable checkpoint. Their intervals of seconds describe the balance for that smaller payload; selecting a recovery cadence requires including optimizer, schedule, and input-stream state. Small-state jobs rarely run on a cluster this large, and their actual storage bandwidth and complete recovery state determine whether saves every few minutes or every epoch have negligible overhead. The large-state rows show why interval selection becomes consequential when each save moves hundreds of gigabytes or terabytes.

Table 7.6: Checkpoint Overhead by Model Type: The Archetype A row is the 175B running example at 14 bytes per parameter (weights, FP32 master weights, and Adam moments); the 20B, recommender, and vision transformer rows are representative scenario sizes, and the 4 TB embedding-heavy recommender is a smaller scenario than the 10 TB+ embedding tables of Archetype B (DLRM at Scale) in table 7.5; the BERT-Large and ResNet-50 rows count FP32 weights only and omit optimizer state, so their write times understate a full resumable checkpoint. Larger models require longer save times but also benefit from longer optimal checkpoint intervals. The percentage column reports the save-side overhead f_{ckpt} of equation 7.9 with no extra pause, not total Young-Daly waste; at the optimum, expected rework adds a comparable term.

Model type	Checkpoint size	Write time (100 GB/s)	Optimal interval (5 h GPU-only MTBF)	Save overhead
Archetype A (GPT-4/Llama-3)	2.45 TB	24.5 s	15.7 min	2.6%

Model type	Checkpoint size	Write time (100 GB/s)	Optimal interval (5 h GPU-only MTBF)	Save overhead
20B dense transformer class	240 GB	2.4 s	4.9 min	0.8%
BERT-Large	1.4 GB	0.014 s	22 s	0.06%
Embedding-heavy recommender	4 TB	40 s	20 min	3.3%
ResNet-50	102.2 MB	0.001 s	6 s	0.02%
Vision transformer class	1.2 GB	0.012 s	21 s	0.06%

These per-model overheads assume the write path is uncontended. When thousands of workers save at the same step, their combined burst becomes the checkpoint storm that section 4.6 defines and sizes; for recovery, the burst matters because a stalled write delays the generation that a failure would fall back on. Staging hides the write from training, yet a snapshot held only on local staging does not survive loss of its node.

7.8.4 Checkpoint storage and recovery

Writing a checkpoint to node-local NVMe makes the save fast, but if the node itself fails, that local snapshot vanishes along with the hardware it was meant to recover. Writing only to a durable remote filesystem or object store survives the node failure, but pulling multi-terabyte state back across shared storage links slows every recovery. Recovery mechanisms therefore pair tiers from the storage hierarchy. Local NVMe provides fast staging and survivor reads, distributed filesystems provide durability across node failures, and object storage provides long-term retention. Together, the tiers let the saved state both outlast the fault and reload within the recovery budget.

↩ Tiered checkpoint storage designs and asynchronous write pipelines are established in section 4.6.

For fault tolerance, checkpoint placement determines which failures the saved state survives, while the restore path determines how quickly that surviving state becomes useful. On the 128-node Archetype A reference cluster, each node holds about 19.1 GB of the 2.45 TB checkpoint for the 175B running example. Reading that share from local NVMe at 5 GB/s takes about 3.8 s, with every node reading in parallel, while pulling the full checkpoint back through shared storage at 100 GB/s aggregate takes about 24.5 s. The local path is faster because it multiplies drive bandwidth by node count, but a sole copy on a lost node cannot restore that node's shard. The replacement node must read its share from a copy outside the failed domain, such as a distributed filesystem or object store, whose recovery time is set by sustained aggregate reads, network capacity, and deserialization. A sound placement therefore keeps a fast local copy for the survivors and a durable remote copy for whatever the failure took away, and these rates must still be measured under the job's own load.

7.8.5 Checkpoint coordination patterns

If workers in a sharded training job save their local shards at different steps without coordinating, rolling one failed worker back to its last save invalidates every peer that already consumed gradients or activations from its later steps. Rollbacks then cascade backward through the run in the classic domino effect (Elnozahy et al. 2002). Avoiding that cascade requires choosing how workers coordinate across two orthogonal dimensions. Workers can vary in whether they synchronize to capture a consistent cut, a set of per-worker states that could all have existed at the same moment (**coordinated** versus **uncoordinated checkpointing**). They can also vary in whether they funnel state through one writer or persist shards in parallel (centralized versus distributed checkpointing). Because distributed ML workers exchange state at every step, production training relies on **coordinated checkpointing**. In this scheme, workers synchronize at step boundaries or through an explicit two-phase commit protocol, so that all saved shards belong to a single globally consistent recovery line,

the consistent cut that every worker rolls back to together. Once coordination guarantees consistency, the remaining architectural decision is how those shards reach persistent storage: whether workers funnel their state through a single coordinator or write their partitions directly in parallel.

7.8.5.1 Centralized checkpointing

In centralized checkpointing, workers send their state to a coordinator process that assembles and writes the complete checkpoint. This approach simplifies checkpoint management and can produce self-contained checkpoint files, but all state crosses the coordinator's network links, and coordinator failure interrupts the save. A streaming writer can avoid buffering the entire checkpoint, although it cannot remove the ingress bottleneck. That bottleneck is structural. The coordinator receives through one node's network links, so its write time grows with total state volume, whereas independent writers each move only their own share of the state through their own links. The gap between the two patterns therefore widens in proportion to the number of writers. For the 175B running example, the 128 nodes of the reference cluster would funnel 2.45 TB through one node's links instead of 19.1 GB through each. The pattern is viable while that one path meets the checkpoint budget; state volume and sustained bandwidth determine the limit more directly than a fixed worker count. Centralized checkpointing is therefore attractive when operational simplicity matters more than aggregate bandwidth, whereas distributed and sharded approaches become necessary once the checkpoint itself is a fleet-scale object.

7.8.5.2 Distributed checkpointing

Distributed checkpointing eliminates this coordinator funnel by having each worker persist its own shard directly. Figure 7.13 contrasts the dataflow topologies. The centralized pattern concentrates traffic at the coordinator's ingress, while the distributed pattern writes in parallel to local NVMe and drains the shards to remote object storage in the background. With fixed per-worker shard sizes and sufficient independent staging bandwidth, adding workers need not lengthen the local write pause. Shared drain capacity and completion coordination can still limit scaling. A coordinator signals when to checkpoint and confirms completion, but state flows directly from workers to storage without aggregation.

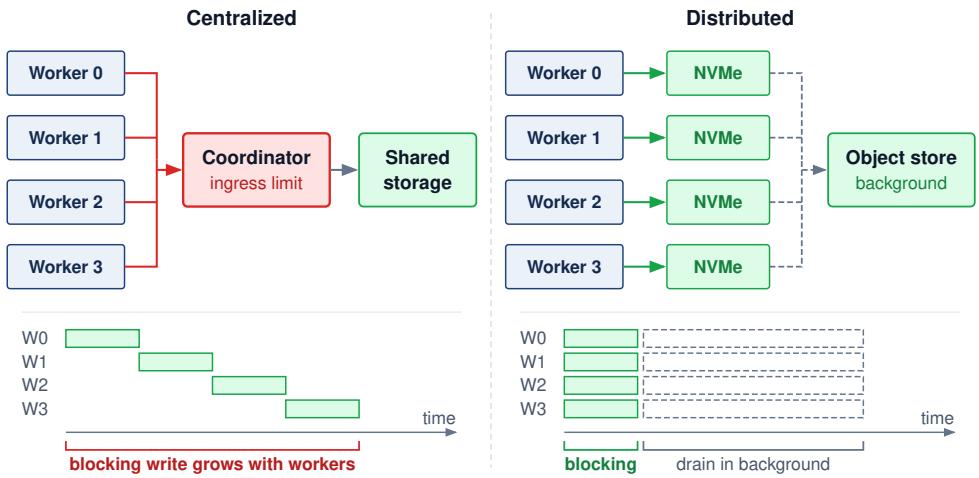

Figure 7.13: Distributed Checkpoint Architecture: Centralized vs. sharded checkpoint writes. In the centralized pattern (left), every worker funnels its shard through one coordinator, whose ingress bandwidth and buffering constrain the write path and whose failure halts the save. In the distributed pattern (right), each worker writes its own shard to local NVMe in parallel and a background process drains the shards to remote object storage, so the local blocking write time can remain roughly constant with fixed shard sizes and sufficient staging bandwidth; shared drain capacity still limits completion.

A synchronous version of the coordination protocol proceeds in six steps:

1. Coordinator broadcasts checkpoint request with checkpoint ID

2. Each worker reaches a consistent state (barrier synchronization)
3. Each worker writes an immutable shard to a generation-specific staging prefix and records its byte length and cryptographic digest
4. Each worker confirms durable write completion to the coordinator
5. The coordinator writes an immutable manifest after all confirmations; the manifest names the training step, topology, model and optimizer formats, data-loader position, random state, and the length and digest of every shard, then publishes a commit marker or atomic pointer to that manifest
6. Coordinator broadcasts checkpoint complete, training resumes

The idealized protocol assumes step 2 completes quickly. The barrier, however, waits for the slowest of thousands of arrivals, so a single slow worker can make it dominate the checkpoint pause even when the storage path has ample bandwidth. Measuring arrival skew separately from copying and writing reveals whether coordination or data movement limits the save.

The manifest, not the presence of a directory or a plausible set of shard names, defines a committed checkpoint generation. Restore accepts only a published manifest whose shards all match their recorded lengths and digests; staged shards without that final publication remain unreachable and can be garbage collected. This rule prevents a partially overwritten `latest` directory from masquerading as recoverable state and makes generation rollback explicit. PyTorch Distributed Checkpoint follows this ordering, with storage writers persisting each rank's data before the coordinator's `finish()` operation writes the metadata and marks the checkpoint successful, while the storage backend supplies the publication semantics (PyTorch Contributors 2026a).

Steps 3 through 5 are two-phase commit,[14] the classic distributed-systems protocol for making one decision atomic across many participants (Gray 1978), applied to a checkpoint. Steps 3 and 4 form its prepare phase, in which each worker stages its shard and reports success, and step 5 is its commit phase, in which the coordinator durably records the decision by publishing the manifest. The manifest publication rule inherits the protocol's guarantee and its known weakness. If the coordinator crashes after workers prepare but before it records a decision, the system must not expose those shards as a valid checkpoint, and the prepared participants remain blocked until the original or a replacement coordinator recovers the durable decision. Uncommitted staged shards can be garbage collected only after recovery establishes that no commit decision exists.

[14] **Two-phase commit (2PC)**: Formalized by Gray (1978), 2PC ensures all participants either commit or abort atomically. A timeout can trigger coordinator recovery, but it cannot by itself authorize unilateral rollback.

Atomic publication also does not by itself make the published shards consistent. Synchronous gradient updates naturally create step boundaries where all workers have applied the same updates, whereas asynchronous training and pipeline parallelism require explicit coordination to define a consistent cut. The set of per-worker states a manifest names is a cut, and the cut is consistent when all of those states could have existed together at one moment of a valid execution. In a consistent cut, no shard reflects an update another shard has not yet seen. If rank 0 saves its shard after step 1000 and rank 1 saves its shard after step 1001, rank 1's shard carries an update that rank 0's shard has not applied, so the pair describes a model the job never held and cannot be restored together.

The coordination protocol decides which cuts a manifest can name, and the consistency model the system claims (strict, bounded-asynchronous, or eventual) states whether that consistent cut is fixed at save time or reconstructed at recovery from retained versions. Synchronous training takes the strict model at its natural step boundaries; the weaker models apply when workers advance asynchronously and share no common step at which to save.

> **Systems Perspective 7.4: Checkpoint consistency models**
>
> Strict synchronous: All workers checkpoint at exactly the same training step. Provides strongest consistency but highest overhead from barrier synchronization.
>
> Bounded asynchronous: Workers may progress at different rates while the checkpoint manager tracks which snapshots can form a consistent recovery cut. Bounding step skew alone is insufficient, because the system must retain compatible shard versions or enough replay information to reconstruct that cut. Asynchronous flushing can also preserve a single synchronized

snapshot while its writes complete in the background, avoiding training stalls without relaxing recovery consistency.

Eventual consistency: Workers save independently and recovery attempts to reconcile their versions. A usable design must retain enough compatible versions or replay information to reconstruct a valid state; arbitrary recent shards do not necessarily form a recoverable generation. Reduced coordination at save time therefore transfers work and storage requirements to recovery.

7.8.5.3 Sharded checkpointing

Modern distributed training frameworks partition model state across workers using techniques like ZeRO and FSDP, which shard parameters, gradients, and optimizer state instead of replicating them on every data-parallel worker. In these configurations, no single worker holds complete model state. Each worker holds only its assigned parameter shard plus corresponding optimizer state.

Sharded checkpointing[15] (Rajbhandari et al. 2020) uses this distribution. Each worker writes only its shard, which divides the per-worker write volume by the number of shards. Recovery loads shards and redistributes state to workers based on the recovery configuration.

↩ How ZeRO and FSDP partition training state across data-parallel workers is established in section 6.3.3.

Dividing state across the fleet keeps per-worker saves small even for multi-terabyte models. A 175B parameter model with a 2.45 TB checkpoint distributed across 1,024 workers requires each worker to write only 2.4 GB, achievable in seconds with local NVMe storage. When elastic scaling or hardware replacement changes the worker count between save and restore, shard redistribution must remap the saved partitions to the new worker configuration, which is why frameworks need flexible resharding support to recover across topology changes. Yet possessing a valid checkpoint is only half the recovery path. Sharding distributes the write work and staging controls the shared I/O burst, but the system must still detect the failure, trigger restoration, and coordinate the distributed shards before training can resume. The speed of detection and recovery therefore governs the cost of the interruption.

15 **Sharded checkpointing**: Each worker saves only its local partition rather than gathering state to a single writer, parallelizing I/O across all nodes. The trade-off is that recovery requires every shard to be available and consistent, making the checkpoint protocol more complex. Losing one storage device invalidates the generation only if a required shard cannot be recovered from another copy or sufficient storage redundancy.

Checkpoint 7.4: Reasoning about checkpoint overhead

Verify your understanding of how failure rate and checkpoint cost set the optimal checkpoint interval:

- ☐ **Young-Daly shift**: use the Young-Daly formula to estimate how the optimal interval T^*_{ckpt} shifts when a cluster doubles in size and its system MTBF halves, with T_{write} held fixed. State whether the interval lengthens or shortens and by about what factor.
- ☐ **Two-sided overhead**: explain why writing checkpoints too frequently and writing them too rarely both raise checkpoint overhead, including the two costs T^*_{ckpt} balances.
- ☐ **Faster writes**: determine how an order-of-magnitude reduction in T_{write} from sharded checkpointing changes the optimal interval and expected lost work per failure.
- ☐ **Square-root scaling**: explain why the Young-Daly interval depends on the square root of MTBF rather than scaling linearly with it.
- ☐ **Overhead at scale**: explain why the optimized checkpoint overhead still rises when a cluster doubles in size with write time held fixed, even though the interval adapts, and by about what factor.

7.9 Failure Detection and Recovery

A GPU silently hangs, dropping its utilization to zero while its peer GPUs wait indefinitely at an AllReduce barrier. On the chapter's 10,000-GPU cluster, every minute the system takes to notice this hung rank and reboot the node costs about $330 in idle compute. Checkpoints preserve state, but the recovery process itself determines how much compute a failure wastes. In the accounting of the Fleet Law (equation 1.3), that waiting belongs to the synchronization term $T_{\text{sync}}(N)$. Every surviving rank idles at the barrier until the group is whole again, so detection, restart, and reload

add directly to step time averaged over the run. The goodput accounting of section 9.2.3 charges the same time to runtime goodput, the share of a job's held allocation spent on steps a checkpoint retains.

Equation 7.10 decomposes recovery time into four primary components:

$$T_{\text{recovery}} = T_{\text{detect}} + T_{\text{restart}} + T_{\text{load}} + T_{\text{warmup}} \tag{7.10}$$

where T_{detect} is the time between the physical hardware fault and system classification as a failure. The term T_{restart} is the latency required for the job scheduler to allocate resources and launch replacement processes. The term T_{load} is the I/O duration to read checkpoint state from distributed storage into GPU memory. The term T_{warmup} is the interval needed to refill data pipelines, compile just-in-time (JIT) kernels, and stabilize training throughput. Each component presents distinct optimization opportunities, and the dominant term varies by cluster configuration. Understanding this decomposition enables targeted investment in the bottleneck rather than uniform improvement across all components.

7.9.1 Failure detection mechanisms

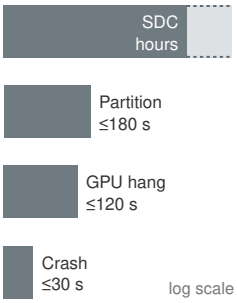

Detection latency climbs from seconds (crash) to hours (silent corruption).

To the rest of the cluster, a dead rank and a rank busy flushing a multi-gigabyte checkpoint shard both look silent. Every failure detector must therefore balance two expensive mistakes: setting timeouts too short kills healthy workers during transient network or storage stalls, while setting them too long leaves thousands of peer GPUs idling at a collective barrier. Heartbeat monitoring is the first detection layer and usually sets the floor on T_{detect}. Each GPU rank periodically signals liveness to a coordinator, and if no heartbeat arrives within the timeout period T_{timeout}, the coordinator declares the worker failed.

For a heartbeat interval of H seconds and network-delay standard deviation σ_d in seconds, equation 7.11 defines the timeout heuristic:

$$T_{\text{timeout}} = H + k\sigma_d \tag{7.11}$$

The first term waits out one full heartbeat interval, because a healthy worker may have just sent its last beat, and the second adds a margin for heartbeats that the network delays longer than usual. Here, k is a dimensionless safety multiplier; values from three to five illustrate a tuning range rather than a guaranteed false positive rate. The heuristic treats the normal heartbeat delivery delay as small relative to the heartbeat interval. A deployment with substantial baseline delay or a heavy delay tail must account for those measurements as well as σ_d. Tuning T_{timeout} too tightly causes legitimate checkpoint writes or long collectives to trigger false failure declarations; tuning it too loosely extends the window during which other ranks wait for the missing contribution.

Collective communication timeouts provide a second detection layer. During synchronous training, a collective required by the next step remains an incomplete critical-path dependency if a required rank disappears. AllReduce needs every rank's contribution before it can produce the reduced result; the progress and failure behavior of other collectives depends on the operation and runtime. Transport and process-group timeouts bound how long the job waits before declaring the operation failed and entering recovery. NCCL[16] provides configurable transport and reliability, availability, and serviceability (RAS) timeout parameters for this purpose, while higher-level frameworks may impose process-group operation timeouts. These settings are often set conservatively to avoid crashing jobs during legitimate periods of slow communication, which unfortunately extends T_{detect}.

Container orchestration health checks provide a third layer. Kubernetes, the container orchestration system, offers liveness probes (verifying that a process still responds and makes progress) and readiness probes (verifying that a process is ready to handle requests). Slurm, the batch scheduler common on HPC-style training clusters, instead runs a periodic node health-check program that can drain a node whose GPUs, links, or file systems fail their tests. Both operate independently of the training framework, catching failures that application-level heartbeats might miss; a liveness probe that tests forward progress rather than mere process existence can catch a process that is alive but deadlocked.

Loss spike detection exposes an anomaly, not its cause. Figure 7.14 contrasts an idealized healthy loss trend with gradient explosion, a premature plateau, and NaN divergence. Real stochastic

16 **NCCL timeout**: NVIDIA's collective communication library exposes transport and RAS timeout controls such as `NCCL_IB_TIMEOUT` and `NCCL_RAS_TIMEOUT_FACTOR` (NVIDIA 2026c). Higher-level launch and distributed-training frameworks can add their own restart and operation-timeout behavior (PyTorch Contributors 2026c). Aggressive timeout settings accelerate failure detection but risk false positives during legitimate long collective operations on large models, forcing a trade-off between detection latency and training stability that has no universal optimum.

training need not decrease monotonically, so a transient rise alone does not establish failure, and nonfinite values need not crash the process. Silent hardware corruption can produce an abrupt loss or gradient-norm spike, but so can a bad batch, an optimizer instability, a data-ordering defect, or a distributed-state mismatch. Robust systems therefore pause before the suspect update commits, preserve the input and worker evidence, and replay the step on known-good hardware. The system drains hardware only after that evidence implicates the device; otherwise, it repairs the data, optimizer, or synchronization path that produced the same symptom.

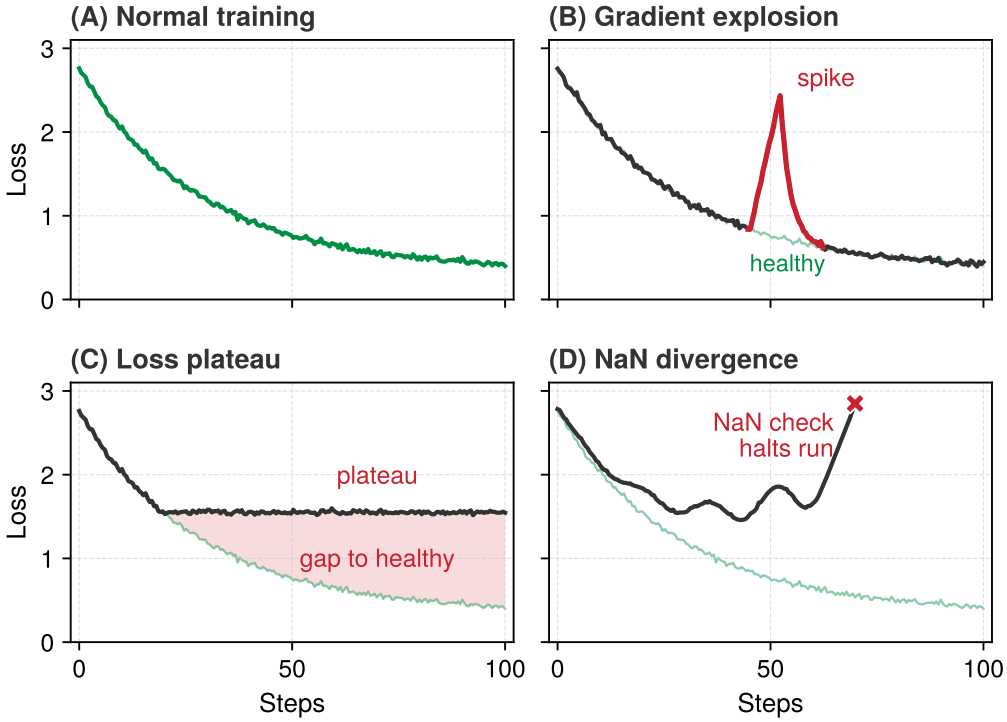

Figure 7.14: Training Failure Signatures: Common patterns of model loss over time that indicate underlying system or algorithmic failures. (A) Normal training. (B) Gradient explosion. (C) Loss plateau. (D) NaN divergence, ending where a nonfinite-value check halts the run; without such a check the process can keep running on NaN values rather than crash. Identifying these signatures automatically is essential for rapid recovery in autonomous training pipelines. The curves are synthetic illustrations of each pattern, not logs from a training run.

Table 7.7 pairs each signature, plus the silent hang that a loss curve cannot show, with a candidate cause and a first response. The pairing guides investigation without treating the curve as a unique diagnosis, and the operational question that follows is how long detection actually takes once these monitoring layers interact.

Table 7.7: Training Failure Response Map: Training-loss curves are operational evidence only when they lead to different recovery actions. The response map ties each signature to the failure class most likely to explain it and to the first action operators should take.

Signature	Likely interpretation	Recovery response
Spike followed by recovery	Transient input or numerical anomaly; the curve alone does not identify the cause	Preserve the suspect batch and step evidence, then replay before deciding whether to continue.
Spike followed by plateau	Optimizer instability, corrupted state, or a persistent data-path defect	Pause, replay from a known-good checkpoint, and compare input, optimizer, and worker evidence.
Gradual divergence	Data drift, numerical instability, corrupted state, or worker desynchronization	Compare per-rank telemetry and evaluation evidence; isolate hardware only when replay implicates it.

Signature	Likely interpretation	Recovery response
Hang without an error	Collective deadlock or crashed worker blocking synchronization	Use timeout detection and restart or reconstruct the worker group.

Systems Perspective 7.5: Realistic failure detection latencies

Failure detection includes more than the configured heartbeat timeout when the policy requires verification before recovery. The ranges in table 7.8 illustrate planning scenarios, not universal detection guarantees. Aggressive timeouts risk false positives (killing healthy-but-slow workers), while conservative timeouts delay real failure detection. A multi-stage policy can use an initial timeout to trigger investigation and a later threshold to trigger recovery. Its actual latency must be measured under the workload and failure conditions it is intended to handle.

Table 7.8: Illustrative Failure Detection Latencies: Planning ranges for detecting process crashes, GPU hangs, network partitions, and silent data corruption, with the dominant cause of latency for each category, showing why static heartbeat timeouts alone are insufficient. The ranges are illustrative assumptions, not measurements; a deployment must measure its own.

Failure type	Planning range	Why
Process crash	5–30 seconds	Heartbeat timeout + verification retries
GPU hang	30–120 seconds	Must distinguish from legitimately slow kernel
Network partition	60–180 seconds	Must distinguish from temporary congestion
Silent data corruption	Minutes to hours	Requires statistical anomaly detection

7.9.2 Recovery procedures

Once a failure is classified, a cold-restart recovery path executes a coordinated sequence to restore consistency. The warm-restart alternative developed next preserves surviving processes when their state can be trusted.

1. Job termination: A `SIGTERM` is broadcast to all surviving workers. This cold-restart path tears down the affected training group so that no survivor continues using the old membership or partially completed collective state.
2. Resource reclamation: The scheduler marks the failed node as "draining" to prevent immediate rescheduling and requests a replacement from the spare pool.
3. Job restart: New containers are launched (from cache if available), and the training binary is re-initialized on all nodes.
4. Checkpoint loading: Each worker reads the state assigned to it from the distributed filesystem. When membership or parallelism changes, recovery remaps the saved shards to the new partitions.
5. State synchronization: Ranks handshake to establish a new communicator, the object that names the participating set (for example, through `ncclCommInitRank`), and workers verify they are all at the same training step.
6. Training resumption: The data loader restores its recorded position and sampling state, and the training loop resumes from the checkpoint step. A batch index alone suffices only when it reconstructs the same logical input stream.

Automatic recovery systems perform these steps without human intervention. Modern training frameworks integrate with cluster managers to automate parts of the sequence. DeepSpeed documents save/load routines for model and ZeRO optimizer checkpoints, which provide the durable state needed after failure (DeepSpeed Developers 2026a). PyTorch's `torchrun` elastic launch documents failure and membership-change behavior through its rendezvous mechanism,[17] the

[17] **Rendezvous**: From French "rendez-vous" (present yourselves), this coordination protocol lets restarted workers discover one another and establish a new group before training resumes. Recovery waits for the membership policy to be satisfied, so slow startup, delayed communication, and missing workers can extend this phase. Heterogeneity affects those delays but does not by itself imply a proportional latency law.

coordinator-backed protocol by which surviving workers agree on membership and ranks after a restart (PyTorch Contributors 2026c, 2026b).

Recovery validation tests whether the restored job behaves consistently with its recorded state. Parameter shapes, dtypes, and shard identities must match the manifest. A short replay or validation workload then checks loss and gradient statistics under the expected numerical tolerance; different batches need not reproduce the last prefailure loss. Immediate divergence is evidence to investigate the checkpoint, data state, and execution configuration, not proof that the snapshot is corrupt. Falling back to an earlier validated generation is appropriate when the newest state cannot be trusted.

Recovery must also prevent the failed ownership epoch from returning. A coordinator grants each worker-group incarnation a monotonically increasing epoch (a membership counter, unrelated to a training epoch over the dataset) and associates membership with a lease, a grant of ownership that expires unless the holder renews it. Every state-changing write carries that epoch as a fencing token. Before replacement writes proceed, the recovery protocol advances the enforced epoch at each protected resource, which rejects writes carrying an older token (figure 7.15). The old worker then cannot publish a checkpoint or update job progress, even if it resumes after a long pause; rendezvous likewise admits only the current membership. A lease alone is insufficient because expiration changes the coordinator's ownership decision, not the old worker's ability to reach storage. Fencing enforces that decision at the resource boundary (Kleppmann 2016).

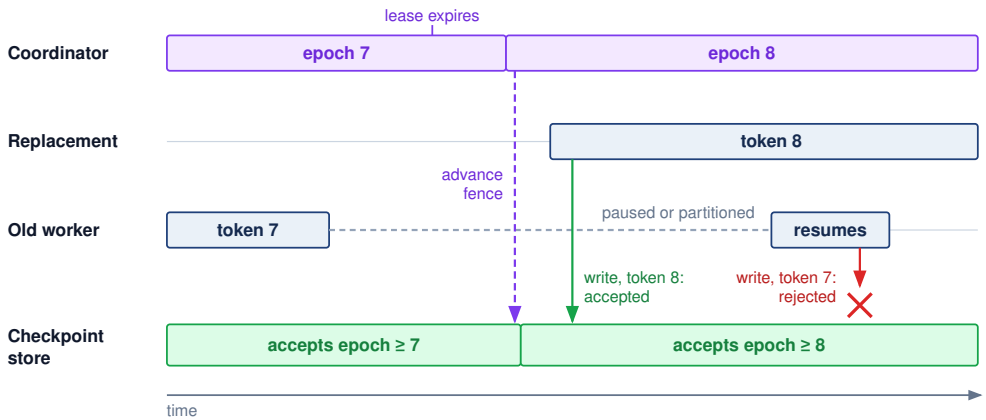

Figure 7.15: Epoch Fencing during Recovery: When the old worker's lease expires, the coordinator starts epoch 8 and advances the epoch the checkpoint store enforces before the replacement begins writing. The replacement's writes carry token 8 and are accepted; when the old worker resumes from its pause or partition, its write still carries token 7 and the store rejects it. The lease decides ownership, and the fence at the store enforces that decision even against a worker that never learned it lost ownership.

Recovery steps such as a checkpoint read or a rendezvous attempt are often retried when they fail, and a careless retry can repeat a side effect or overload a service that is itself recovering. An operation may be retried automatically only when it is idempotent, meaning that repeating it has the same effect as performing it once, or carries a stable operation identifier that lets the receiver deduplicate repeated effects. Each attempt inherits the caller's absolute deadline rather than receiving a fresh timeout, and the system caps attempts with a retry budget shared across the request or recovery tree. Exponential backoff, which multiplies the wait by a fixed factor after each failed attempt, and jitter, which randomizes each wait, spread attempts in time, but they do not prevent overload if every layer retries independently. The budget therefore belongs at one retry-owning layer; exhausted budget returns a failure or invokes degradation instead of amplifying the original incident (Beyer et al. 2016; Nygard 2007).

Napkin Math 7.3: The recovery time budget

Problem: Consider a 175B parameter model training on 1,024 GPUs. The checkpoint size is approximately 2.45 TB (BF16 weights, FP32 master weights, and Adam moments at 14 bytes per parameter). Estimate the cost of a single failure event.

Setup: The recovery budget T_{recovery} has four terms.

- $T_{\textbf{detect}}$: 60 s (conservative heartbeat timeout with verification retries).
- $T_{\textbf{restart}}$: 3 min (scheduler queue time + container launch + Python import overhead + NCCL initialization).
- $T_{\textbf{load}}$: 24.5 s (this is the aggregate-shared-storage read of the full 2.45 TB checkpoint at 100 GB/s aggregate). With sharded local-NVMe reads at 5 GB/s, one GPU's 2.4 GB shard alone would load in under a second. The GPUs of a node share its drive, however, so the node's 19.1 GB share takes about 3.8 s. That local read is available only on nodes that survived with their local copy.
- $T_{\textbf{warmup}}$: 2 min (JIT kernel compilation, data pipeline buffer fill, TCP connection re-establishment).

Total: $T_{\text{recovery}} = T_{\text{detect}} + T_{\text{restart}} + T_{\text{load}} + T_{\text{warmup}} \approx 6.4$ min per failure event.

Impact: For this recovery path, T_{recovery} is the operational MTTR term in the steady-state availability relation (perspective 7.1). It is not the whole cost of a failure. Every failure also discards the work done since the last completed checkpoint, on average half the checkpoint interval (figure 7.12), and the Young-Daly waste of equation 7.8 already counts that rework, so the figures below price only the recovery term.

Recovery time optimization matters more as clusters grow. The GPU-only baseline from table 7.1 compares the two sizes on a common basis and understates both failure counts. On that basis, a 1,024-GPU cluster has MTBF of 49 hours and experiences about 0.49 failures/day, losing about 3 minutes daily, a modest 0.2 percent overhead. However, a 10,000-GPU cluster has GPU-only MTBF of 5 hours and experiences about 4.8 failures/day, losing about 31 minutes daily (2.1 percent overhead). That is equivalent to wasting roughly 5,127 GPU-hours every day (the daily lost minutes multiplied across all 10,000 GPUs). Failures arrive in proportion to the number of GPUs, so when each one costs the same recovery time, recovery overhead grows linearly with cluster size. The optimized checkpoint terms grow only with the square root of cluster size (section 7.8.1), so under that fixed recovery time, recovery eventually becomes the larger loss.

7.9.3 Warm restart vs. cold restart

The recovery procedure described in section 7.9.2 is a cold restart: every process in the affected job group exits, and the group reloads state from a valid checkpoint. This simplifies recovery by avoiding dependence on surviving in-memory state, but it can discard useful state. After one GPU fails in the 1,024-GPU reference job, every surviving worker may retain reusable values; an interrupted update or silent corruption, however, can make some of those values unsafe to keep.

A warm restart preserves the state of surviving workers. When a rank fails, surviving ranks detect the failure but do not exit. They enter a waiting state, preserving loaded model weights and optimizer states in GPU memory. The scheduler replaces only the failed node. The new node joins, receives the current state of its partition from a peer replica where one exists, or otherwise loads it from the latest checkpoint, and the communicator is rebuilt. The replacement and survivors must agree on the same committed training step before resuming. A shard restored from an older checkpoint cannot be combined with survivor shards that have advanced beyond it; the runtime must reconstruct the missing current state or roll the group back to a common recovery point. Training resumes only after that consistency check.

When the missing state can be reconstructed at the survivors' current step, warm restart avoids a full-cluster reload and repeated warmup. For the 1,024 GPUs example, that fast path avoids rereading the full 2.45 TB checkpoint through shared storage, whose cold-path T_{load} is 24.5 s. The replacement

still needs state transfer and initialization, and all ranks wait until it is ready. If consistency instead requires a rollback, the recovery budget must include the corresponding survivor-state restoration and lost work.

The trade-off is software complexity. Survivor-preserving warm restarts require specialized application and runtime support for dynamic membership changes without leaking CUDA memory, corrupting shared state, or deadlocking during communicator reconstruction. Common training frameworks restart the worker group or reload a checkpoint rather than preserve survivors (section 7.10.1.2). A deployment that implements a warm-restart fast path still needs cold restart as a safety net. The cold-restart range in table 7.9 brackets the 6.4 min budget computed above. The warm-restart range is a planning range for the fast path, which applies only when the replacement's current state can be reconstructed and avoids the full-cluster reload and warmup; its actual duration depends on how quickly the replacement receives its state and the communicator is rebuilt.

Table 7.9: Cold vs. Warm Restart Trade-Offs: Cold restart kills every process and reloads from disk; warm restart preserves surviving workers and replaces only the failed node. Warm restart cuts recovery from minutes to tens of seconds for the typical single-node failure but introduces software complexity around dynamic membership management. Mature deployments can run warm restart as the fast path and cold restart as the safety net. Recovery times are illustrative planning ranges, not measurements.

Aspect	Cold restart	Warm restart
Recovery time	4–10 minutes (full reload)	30–90 seconds (single node reload)
State guarantee	Clean: all state from checkpoint	Assumes surviving state is valid
Implementation	Simple: kill all, reload all	Complex: dynamic membership management
Failure during recovery	Retry cold restart	Fall back to cold restart
Best for	Correlated failures, SDC events	Single-node failures, GPU errors

7.9.4 Recovery automation pipeline

At the scale of 10,000+ GPUs, failures occur multiple times per day, so human operators cannot intervene manually on each incident. Recovery must be an autonomous control loop managed by the cluster control plane. Published reports from large training fleets describe multi-stage automation that classifies failures and selects the minimum viable remediation (Dubey et al. 2024; Kokolis et al. 2025; Zu et al. 2024).

The pipeline operates in four ordered stages:

1. Health monitoring daemon: Continuously scrapes GPU telemetry (ECC error counters, temperature, fan speed, NVLink status) alongside application metrics like training loss and step throughput, often from a sidecar container, a helper process deployed beside the training container.
2. Failure classifier: Determines whether a signal indicates a fatal error (for example, NVIDIA Xid error 48: double-bit ECC error), a transient stall (for example, temporary network congestion), or a performance degradation (for example, thermal throttling).
3. Action selector: Chooses the cheapest response consistent with the affected state and communication group. A process hang on one rank usually requires restarting the whole communication group rather than one container, because every surviving rank still holds a communicator that includes the hung rank. A confirmed hardware fault requires draining the node and drawing a replacement from the spare pool, which is slower and consumes standby capacity. A network partition can justify a bounded wait that preserves in-memory state, but if the wait expires, replacement must fence the old ownership before accepting new writes.
4. Validation stage: Runs a recorded "canary batch" or validation workload under the expected numerical tolerance. Unexpected loss or gradient behavior triggers investigation of restored state and execution configuration, with fallback to an earlier validated snapshot when the new state cannot be trusted.

This automation removes the human page-diagnose-repair cycle from the recovery path for the common failure types, which shortens mean time to recovery. The classification stage matters most.

Treating every failure as a cold restart wastes compute on transient issues, while treating a hardware failure as transient allows corrupted computation to continue. The Llama 3 pretraining run shows what such a pipeline buys at fleet scale.

War Story 7.2: The training run that failed every few hours (2024)

Context: Meta reports training Llama 3 405B on 16,384 H100 GPUs during a 54-day snapshot of pretraining (Dubey et al. 2024).

Mechanism: The snapshot recorded 419 unexpected interruptions, about one every 3.1 hours, with the root causes shown in figure 7.2.

Impact: Meta nevertheless reported more than 90 percent effective training time, measured as useful training time divided by elapsed time. Significant manual intervention was required only three times.

Fix: Meta reduced job startup and checkpointing time and developed automated diagnosis and recovery tools. Communication traces helped locate stalled collectives and slow workers before restarting useful work.

Systems lesson: At fleet scale, reliability is a property of the recovery system rather than of the hardware. Effective training time stayed high because automation, not operators, handled nearly every interruption.

7.9.5 Stragglers versus fail-stop failures

Not every recovery decision starts from a crashed process. When one GPU in a synchronous job slows down without dying, every healthy peer still waits at the AllReduce barrier, forcing the recovery pipeline to decide whether the slow rank is a temporary straggler or an incipient failure (perspective 7.6).

Systems Perspective 7.6: Straggler or failure

A straggler, the slow rank of a synchronous step (section 6.3.1.5), bounds the throughput of every peer that waits on it, so one node running 10 percent slow can cost thousands of nodes 10 percent of their capacity. Fault tolerance must decide what that slowness means. A straggler can still produce correct results, so slowness alone establishes neither incorrect values nor a failed process, and its cause is often system jitter from background OS processes, network congestion, or thermal throttling rather than bad hardware. The decision is between mitigation, which keeps the job running and removes the slowdown, and recovery, which treats the rank as failed and pays for a restart.

Under the simplified no-overlap straggler model, the slowest rank sets the step time of the entire cluster because AllReduce cannot complete until every participant has submitted its gradients:

$$T_{\text{step,straggler}}(N) = \max(T_{\text{rank}_0}, T_{\text{rank}_1}, \ldots, T_{\text{rank}_{N-1}}) + T_{\text{comm}}(N)$$

Many stragglers are gray failures, the silent degradations named in the failure taxonomy, seen from inside a synchronous job, and none of their common causes raises an explicit error. Thermal throttling reduces clock speed, degrading interconnect cables increase communication latency, OS background processes (memory scrubbing, log rotation) consume CPU cycles, and data loading from a congested network filesystem introduces variable I/O delays. A slow worker can continue meeting a liveness timeout while reducing training throughput, or become slow enough to trigger the same timeout as a failed worker.

Progress measurements are therefore needed alongside liveness checks. A common detector records each rank's step time, or the time its contribution reaches each collective, and compares it with the median across ranks, so a rank that is persistently slower than its peers stands out even while it answers every heartbeat. It flags a rank once its time exceeds the median by a set margin over several consecutive steps. A tight margin or a short window catches a degrading device sooner

but also flags ranks slowed by transient jitter, while a loose margin leaves the whole job running at the slow rank's pace for longer.

Timeout policy governs where the system draws the line between *mitigation* (redistributing work or replacing the slow node) and *recovery* (checkpoint-restart). Aggressive timeouts treat transient stragglers as fail-stop crashes, causing unnecessary job restarts that waste more compute than the slowdown itself, whereas conservative timeouts leave thousands of healthy GPUs waiting on a degraded node that will never speed up. A slow rank that stays in the job lowers program goodput, the ratio of a step's ideal execution time to its actual time, on every step, while replacing it spends runtime goodput on the interruption. The goodput accounting of section 9.2.3 records these as separate losses. The replacement decision therefore weighs one loss against the other.

Straggler mitigation strategies change either the work assignment or the execution resources. Dynamic load balancing moves work away from slow workers while preserving the intended update, whereas bounded staleness allows delayed gradients under an algorithm that tolerates them; its convergence cost depends on the workload and delay bound. Backup workers replicate selected tasks and accept one result per logical contribution, trading extra GPU cycles for a shorter wait without double-counting gradients. At the infrastructure boundary, proactive replacement uses hardware telemetry, such as rising operating temperatures or growing correctable-ECC error counts, to identify suspect components and replace them when the expected benefit exceeds migration and recovery costs. A simple cost calculation shows why a severe straggler is often cheaper to replace than to tolerate.

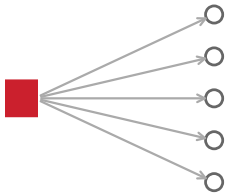

One straggler stalls every healthy peer at the barrier.

Napkin Math 7.4: The cost of a straggler

Consider a 1,024-GPU cluster where a normal training iteration takes 1 second. A single GPU enters a thermally throttled state, clocking down to 50 percent speed, and now takes 2 seconds to complete its computation.

Because AllReduce cannot complete until every rank has submitted its gradients, the other 1,023 healthy GPUs sit idle waiting for the straggler.

Impact:

- Normal step time: 1 second
- Straggler step time: 2 seconds
- Effective cluster speed: 1 step/2s = 0.5 steps/sec

A single failing device (0.1 percent of the hardware) has reduced the throughput of the entire cluster by 50 percent. At \$2/GPU-hour, this straggler wastes \$1,024/hour in idle compute.

Systems insight: Replacement is worthwhile only when the expected time saved exceeds recovery and replay costs. Suppose the slowdown is expected to persist for the next K useful updates, whose step times are T_{slow} on the degraded allocation and T_{healthy} after replacement. The simplified break-even condition is then $K(T_{\text{slow}} - T_{\text{healthy}}) > T_{\text{restart}} + T_{\text{replay}}$. The right side includes waiting for capacity and restoring state. A short-lived slowdown may clear before replacement pays off; false-positive detection and uncertain persistence also affect the decision. MTBF describes failure recurrence, not how long this slowdown will last.

Replacing a slow node while maintaining the original GPU count requires available capacity. An $N+1$ redundancy model reserves a hot spare for the isolation and promotion stages of figure 7.10, reducing allocation delay. Promotion still requires compatible resources, valid state, and communication-group reconstruction; an idle device alone does not guarantee recovery within seconds. A cold spare adds startup work, while a job without a reserved spare depends on the scheduler finding capacity or on repair. The hot spare buys shorter waiting time with devices that sit idle while the fleet is healthy. Fixed-size checkpoint-restart recovery restores the original resource count. Elastic recovery relaxes that requirement by allowing the job to continue with fewer workers when the remaining allocation can support it.

Checkpoint 7.5: Tuning detection and recovery

Verify your understanding of how detection latency and recovery strategy trade against one another:

- ☐ **Timeout tuning**: given the heartbeat timeout rule $T_{\text{timeout}} = H + k\sigma_d$, explain what failure mode is traded away when k is lowered from 5 to 3, and which failure mode worsens.
- ☐ **Warm vs. cold restart**: choose a warm restart or a cold restart for a single GPU failure in a 1,000-GPU job with surviving state intact, and identify two conditions that push recovery back toward the cold path.
- ☐ **Hot spare cost**: explain why a hot spare cuts recovery latency relative to detect-then-restart, including the standing cost the cluster pays to keep that spare available.
- ☐ **Straggler decision**: determine what governs whether to wait out a slowed node or kill it and recover onto healthy hardware.

7.10 Elastic Recovery

Suppose a 1,024-GPU training job loses an 8-GPU node to a hardware fault, but the cluster has no spare nodes available. A fixed-size recovery policy leaves the remaining 1,016 GPUs waiting until replacement capacity becomes available or the allocation is released. Elastic recovery allows the job to resize and continue with less compute, provided its state and parallel decomposition fit the smaller allocation.

Definition 7.2: Elastic training

Elastic training is the capability of a distributed training job to adjust its worker count across membership changes, potentially through a coordinated worker-group restart.

1. **Significance**: It converts a permanent capacity loss into a temporary interruption followed by reduced throughput. A pure data-parallel 1,024-GPU job that loses 8 GPUs retains 99.2 percent of its device count and may resume if the remaining devices can hold the state. Tensor- and pipeline-parallel jobs additionally need a feasible partition and compatible group sizes; retained device count does not guarantee the same fraction of throughput. The recovery policy must either preserve the effective batch size through additional accumulation or validate the optimization changes caused by a smaller batch.
2. **Distinction**: Unlike fixed-size checkpoint-restart recovery, which waits for replacement hardware to restore the original allocation, elastic recovery can restart or resume with a different worker count. The optimization loop may pause while membership, ranks, and state are re-established.
3. **Common pitfall**: A frequent misconception is that elasticity is "automatic." In reality, recovery requires coordinated adaptation: the training framework, the data loader, and the orchestrator must all synchronize their state to ensure that no data samples are lost or double-counted during the resize.

A small example shows the lifecycle in three phases (figure 7.16). Execution begins with a balanced four-node AllReduce ring at full throughput. When Node 3 drops out, the survivors pause at a rendezvous barrier and reload the last consistent checkpoint. The group then contracts into a three-node ring in which each worker takes a one-third larger data shard, resuming productive training at 75 percent throughput if the per-worker rate remains unchanged.

7.10.1 Recovery adaptation mechanisms

Shedding a lost node narrows the active ring instead of idling the whole allocation until replacement hardware arrives, so elastic recovery converts a job termination into a bounded capacity adjustment. When a failure reduces the worker count, the recovery path must adapt several training components before productive work resumes. Each adaptation addresses a specific consistency requirement that, if violated, would corrupt the model or waste computation.

Batch size adjustment comes first. With fewer surviving workers, either each worker must process more samples to preserve the global batch size, or the global batch size must shrink, which in turn may require recalibrating the learning rate to preserve convergence. Changing the global batch size changes the optimization regime as well as the amount of work per step. The linear and square-root learning-rate rules of section 6.4.5 are candidate policies for a resized job, with the worker-count ratio standing in for the batch ratio when per-worker microbatch size and accumulation count stay fixed (Goyal et al. 2017). For a shrink, the square-root rule reduces the learning rate less than linear scaling does; that arithmetic does not establish which policy is more stable, and either policy requires validation for the optimizer, schedule, and workload. Preserving the global batch instead can avoid introducing this change in optimization regime.

Gradient accumulation offers that alternative. Each surviving worker processes additional micro-batches before synchronization to maintain the original effective batch size. If worker count drops by half, doubling the accumulation steps preserves the number of examples per update and can retain the existing learning-rate schedule. Each worker then performs twice as much micro-batch computation per update, although communication and overlap determine the wall-time increase. Equal batch size does not imply bitwise-identical updates or identical behavior from operations that depend on micro-batch statistics.

Data loader redistribution must agree with the recovered training step and new membership. When workers are lost, the loader reassigns the remaining logical input stream so that the resize itself neither omits samples nor counts them twice in an update. Replaying work discarded by checkpoint rollback is expected; duplicating contributions within the resumed stream is a different error. The saved data position and sampling state distinguish those cases.

State resharding adds further complexity when using sharded model parallelism (ZeRO/FSDP). When parameters and optimizer states are sharded once across the entire job without replication, losing a worker removes the only live in-memory copy of its shard, so the job must reload a complete durable checkpoint generation before repartitioning the state across the smaller worker set. When

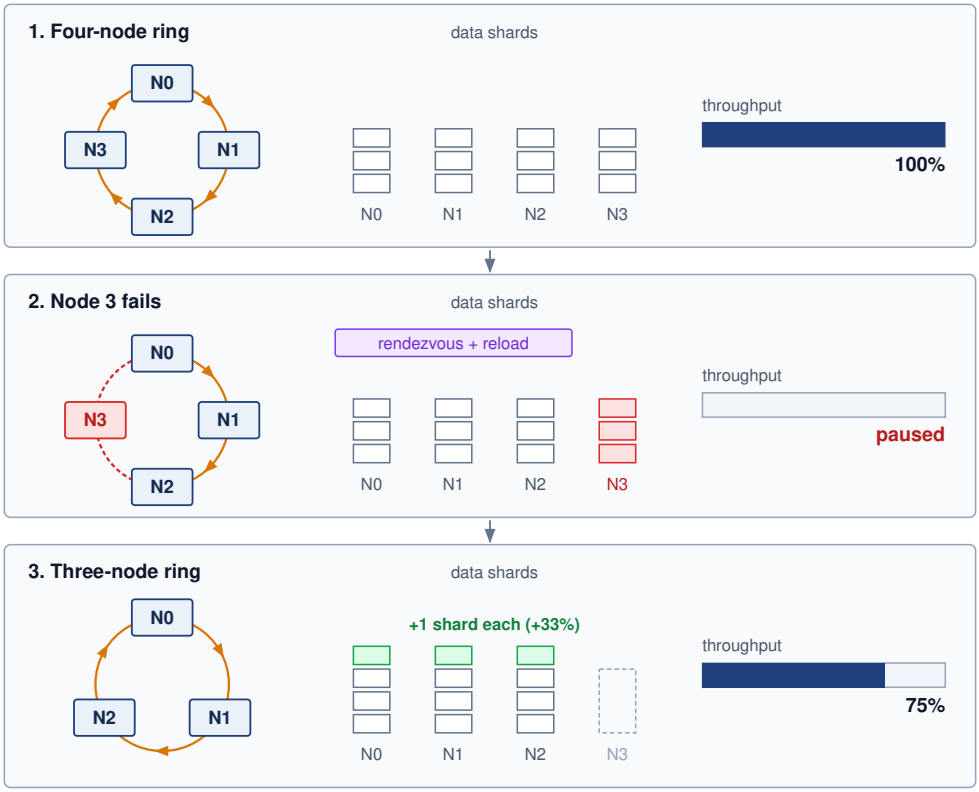

Figure 7.16: Elastic Training Recovery: Read the three phases top to bottom. Rather than restoring the original worker count before resuming, elastic training pauses when a node fails, reloads the last consistent checkpoint, re-forms the ring over the $N-1$ survivors, and redistributes the failed node's data shards among them. The job resumes at reduced throughput (75 percent here, assuming an unchanged per-worker rate), turning a hard failure into a temporary slowdown.

the job instead uses hybrid sharding that replicates each shard across two or more data-parallel groups, a surviving replica group still holds a live copy in GPU memory and can supply the missing shard over the network without reading from storage. Only after the missing values are recovered from a peer replica or a durable checkpoint can the runtime redistribute shard ownership across the surviving workers. Changing shard ownership alone cannot recover missing values.

7.10.1.1 Spot instance preemption as a failure class

Hardware faults are not the only event that removes workers from a running job. Cloud providers often price preemptible ("spot") instances below on-demand capacity, with the caveat that they can be reclaimed with little warning. Once workers vanish, recovery faces the same missing-state and membership problems as a hardware failure. An advance notice, when available, can allow a coordinated save or drain before departure and therefore changes the recovery opportunity.

Both fixed-size and elastic jobs can recover from preemption using checkpoints, but a fixed-size job must regain its required allocation before restarting. An elastic job can instead resume at reduced scale if enough capacity remains and the training invariants survive the resize. Recovery still pays for detection, state restoration, and group reconstruction; repeated reclamations can erase the apparent price advantage. At the reference fleet prices of \$0.70/GPU-hour for spot capacity against \$2/GPU-hour for on-demand capacity, the savings can outweigh the lost work and resizing delays; section 8.8.2 develops the full spot-versus-on-demand cost model.

Lighthouse 7.2: Archetype C (Federated MobileNet): Dropout as the normal case

Archetype C (Federated MobileNet) makes worker departure an expected part of training. Training proceeds in rounds, in each of which a coordinator sends the current model to a set of clients and combines the updates they return. The phone-class clients train while charging, idle, and connected, so a client that disappears mid-round is an ordinary event rather than a fault to repair. The coordinator therefore selects more clients than a round needs and closes the round at a deadline. The round tolerates departures as long as enough compatible updates arrive; otherwise the coordinator abandons it (McMahan et al. 2017; Bonawitz et al. 2019). Model-version checks prevent updates computed against a replaced model from entering the current aggregate, the federated counterpart of fencing a stale epoch, while aggregation weights account for the examples contributed under the chosen learning algorithm. Version checks and weighting do not remove participation bias. Recovery must preserve the coordinator state needed to continue that algorithm, including the model, round position, and any persistent optimizer or sampling state, rather than assuming weights alone suffice. Section 11.7.1 develops the round design, including deadlines, participation requirements, and the bias that arises when some kinds of devices report less often than others.

7.10.1.2 Framework support for elastic recovery

Whether worker loss comes from a hardware fault or a spot preemption, training frameworks differ in how many of the resize invariants (group membership, shard ownership, data loader redistribution, and batch-size or learning-rate adjustment) they enforce automatically after a failure. PyTorch Elastic (TorchElastic) uses agents to monitor workers and a rendezvous mechanism to establish group membership (PyTorch Contributors 2026c, 2026b). On a failure or membership change, it stops all surviving workers, forms a new worker group, and starts its workers with new `RANK` and `WORLD_SIZE` assignments. It handles rank reassignment and communication-group reconstruction while leaving batch size, learning rate, and checkpoint semantics to user code. DeepSpeed (Rasley et al. 2020) provides ZeRO-based distributed training and checkpointing building blocks for large models (DeepSpeed Developers 2026a). Recovery reloads from the last durable checkpoint, and Universal Checkpointing addresses portability across some changes in parallelism and topology (DeepSpeed Developers 2026b), though elastic behavior still depends on the surrounding launcher and resource manager.

Ray Train, built on Ray's actor model, provides a checkpoint-oriented recovery path that restarts a worker group after a failure or preemption and resumes from the latest available checkpoint. The training function remains responsible for saving and loading sufficient state (Ray Project 2026). Horovod Elastic similarly builds on Horovod's data-parallel training system (Sergeev and Balso 2018) with an elastic state API for worker additions and removals (Horovod Developers 2026). The API resets ranks and reconstructs communication while leaving convergence-sensitive choices such as learning-rate or batch-size policy to the training script. Across these four frameworks, restarting the worker group is largely automated, while state resharding remains manual, deployment-dependent, or tied to a checkpoint reload (table 7.10).

Table 7.10: Elastic Recovery Framework Comparison: Frameworks differ in how they detect worker failures and how much of the recovery sequence (group reconstruction, state redistribution, batch size adjustment) they automate.

Framework	Failure detection	Automatic recovery	State resharding	Cluster integration
PyTorch Elastic	Rendezvous timeout	Yes	Manual	Kubernetes
DeepSpeed	External	Checkpoint-based	Deployment-dependent	External schedulers
Ray Train	Actor supervision	Checkpoint retry path	Checkpoint reload	Ray Cluster
Horovod Elastic	Driver heartbeat	Yes	Manual	Slurm, Kubernetes

7.10.2 Model-specific training fault tolerance

Restoring model weights alone does not guarantee a valid restart. An LLM that reloads its weights but loses its curriculum cursor, the record of where training stands in its planned document order, trains twice on the same documents, and a continuous recommender that rolls back its embedding tables serves stale user preferences. The checkpoint and recovery mechanisms of section 7.8 and section 7.9 must therefore adapt to the workload costs priced in section 7.7.1. The design question is not simply how often to checkpoint, but which state variable would make a restart mathematically or economically wrong. For recommendation workloads, incremental checkpointing, tiered checkpointing, and embedding versioning can reduce the cost of preserving fresh state. Incremental storage can still reconstruct an exact generation from a base and its required deltas; accepting stale or approximate state is a separate choice. For each workload, the state at risk determines what its checkpoint must contain (table 7.11).

Table 7.11: Training Recovery State Invariants: Workloads require different checkpoint contents because they make different state variables quality-bearing. The correct recovery policy starts from the state that would make a restart wrong, then chooses sharding, incrementality, or domain-state capture accordingly.

Workload	State at risk	Checkpoint implication
LLM training	Optimizer state, curriculum position, document position, long-context schedule	Preserve the FP32 optimizer state where precision matters, shard writes with ZeRO/FSDP, move serialization off the critical path, and include the data-schedule position so the optimization trajectory resumes correctly.
Recommendation	Embedding freshness, feature-store version	Preserve fresh state with incremental checkpointing, tiered checkpointing, and embedding versioning; accepting stale or approximate state requires a separate recovery contract.
Vision	Augmentation seeds, shuffling order, batch-normalization statistics, progressive-resizing schedule	Capture the stochastic and normalization state that controls the input stream; weights alone are not enough to reproduce the training trajectory.
Scientific ML	Simulator state, search frontier, explored configurations, validation seeds	Treat the model, simulator, search process, and random state as one consistency unit so recovery does not duplicate exploration or invalidate comparisons.

The common pattern is that recovery correctness is broader than weight restoration. LLMs can reload weights but resume with the wrong curriculum position. Recommendation models that train continuously can reload dense parameters but serve stale embeddings, so a failure that loses recent updates carries a freshness cost as well as a replay cost. The recovery budget must reflect the workload's measured sensitivity to stale state.[18] Vision models have smaller checkpoints, about 1.2 GB for the ViT-Large example, which permit more frequent saves on the same storage path. Yet reproducing their input stream requires the augmentation configuration, random-number generator state, and data position rather than only the initial shuffling seed. A progressive-resizing schedule, which raises image resolution as training proceeds, must also resume at the same stage. Scientific models can reload a neural network while losing the simulator state that made the data meaningful, and an iterative search must also retain which configurations it has evaluated when those results choose its next candidates. In every case the checkpoint boundary must include whatever prior decisions affect future work, wherever that state resides. Once the trained model is deployed to users, however, the operational calculus shifts from protecting *weeks of batch computation* to protecting *milliseconds of real-time latency*.

[18] **RecSys freshness**: Freshness makes recovery delay part of the model-quality budget. Its importance depends on how quickly useful user and item signals change and how strongly predictions depend on them. Comparing quality at different state ages establishes the permitted staleness; that measured limit determines how much checkpoint and replication overhead is justified.

7.11 Serving Fault Tolerance

When users ask a voice assistant to turn off the lights, they will not tolerate a five-minute pause while the inference server reloads from a checkpoint. Serving presents a different challenge, because users expect millisecond-level responsiveness even when backend GPUs crash. The C^3 budget for fault tolerance shifts accordingly. Training pays for coordination in synchronization time so that a failure does not erase completed compute, while serving pays in standing redundant compute and in routing coordination so that a failure does not break a latency contract.

A serving replica is a runnable copy of the model that can accept requests; the routing layer chooses a replica, and a readiness check determines whether that replica can serve successfully. Replicas may share external state stores or keep request state in accelerator memory, so replacing compute does not necessarily replace the state a request needs. Serving fault tolerance asks how those mechanisms behave when replicas, GPUs, or state stores fail under live latency budgets.

7.11.1 Stateless vs. stateful serving

The first serving decision is where request state lives, because that choice determines whether failover is a retry or a state-reconstruction problem. A stateless model worker retains no required per-session state between requests. It obtains the necessary context from the request or an external state store, so another ready worker can process the same request. Because the external store carries the session history, a stateful application can still use stateless model workers, although that store remains a reliability dependency.

The stateless pattern appears when the input itself carries all context: an image classifier processes one image, an object detector processes one frame, a single-turn text classifier processes one snippet, and an embedding service maps one input to one vector. Fault tolerance can then focus on replica health and request routing. Redundant replicas serve requests in parallel, load balancing sends traffic only to healthy replicas, health checks remove failed replicas from rotation, and automatic replacement starts a new replica when capacity falls. When a replica fails, in-flight requests to that replica can be retried elsewhere, because no quality-bearing session state has been lost.

Stateless serving offers the simpler architecture when the application permits it. However, many ML applications inherently require state across requests. Consider a chatbot. A single-turn question-answering system can operate statelessly, processing each question independently. A conversational assistant must preserve its history somewhere, but the model worker can still be stateless if each request carries the complete history or reads it from a durable external store. Worker-local state creates a stronger dependency when later requests rely on information available only at the selected replica.

Stateful serving appears when later work depends on retained session state. For a transformer, the key-value (KV) cache stores attention keys and values computed for earlier tokens so the worker can reuse them while generating subsequent tokens. This cache is derived state. If the authoritative token history and compatible model version remain available, a replacement can recompute it by processing that history again. Losing the history itself is a different failure because cache reconstruction no longer has its required input. Streaming speech recognition similarly maintains context from previous audio, recommendation sessions accumulate user context, and interactive editing maintains document state across edits.

The fault-tolerance consequence is that failure loses quality-bearing state rather than serving capacity alone. KV cache loss requires reprocessing previous turns, session context loss forces users to repeat previous interactions, and accumulated user state loss degrades quality when context is unavailable. The mitigation choice follows the size, update rate, and quality value of that state. Session affinity routes a session to the same replica, state checkpointing periodically saves session state, state replication maintains a standby copy, and graceful degradation keeps the service available at reduced quality if state cannot be recovered. Table 7.12 contrasts how that state boundary reshapes routing, failure impact, and recovery complexity. Whether the architecture is stateless or stateful, maintaining availability under live latency constraints requires deploying redundant capacity.

↪ How the KV cache grows with sequence length and concurrency until it fills accelerator memory is developed in section 10.7.1.

Table 7.12: Stateless vs. Stateful Serving Fault Tolerance: State recovery adds work after failure detection. End-to-end failover also includes probe timeouts, routing, replacement readiness, and request retry; neither worker category implies a fixed latency.

Aspect	Stateless serving	Stateful serving
Request routing	Any replica	Session-affine replica
Failure impact	Retry on another replica	Potential state loss
Recovery complexity	Restart and load weights	Reload state + reconstruct context
Redundancy approach	Active-active replicas	Replicated state + standby
Work after detection	Route or retry to a ready worker	Recover required state, then route or retry

7.11.2 Redundancy and replication

Redundancy buys availability only when replica placement and spare capacity match the failure domain being defended against. Turning multiple model copies into uninterrupted service requires sizing the replica count to the availability target, choosing how live traffic is split across healthy and standby workers, and placing those replicas across independent failure domains.

7.11.2.1 Availability calculations

Serving redundancy is the parallel composition of section 7.2.1 applied to interchangeable replicas, with availability in place of survival probability. For a single replica with availability A_{single} (probability of being operational at any given time), equation 7.12 quantifies how multiple independent replicas achieve higher system availability:

$$A_{\text{system}} = 1 - (1 - A_{\text{single}})^k \tag{7.12}$$

where k is the number of replicas. For a service whose single replica is available 99 percent of the time, redundancy compounds as follows.

For a single replica, availability of 99 percent corresponds to 3.65 days of downtime per year. Adding independent replicas changes the tail of the failure distribution quickly. Two replicas give 99.99 percent, or 52.6 minutes of downtime per year, and three replicas give 99.9999 percent, or 31.5 seconds. The arithmetic assumes independent failures and enough capacity on any surviving replica to meet the service target; shared power, shared networking, and shared software bugs reduce actual availability below these theoretical values.

7.11.2.2 Replication strategies

Redundancy turns days of downtime into seconds.

The availability equation says how much redundancy can help; the replication topology determines how that redundancy behaves under normal execution and sudden hardware loss. Two fundamental designs make opposite bets (figure 7.17). Active-active replication splits ingress traffic evenly across live replicas, one-third each for three, which keeps every device busy but requires spare headroom so the survivors can absorb a failed replica's share. Active-passive replication sends all traffic to a designated primary and keeps a pre-warmed standby whose weights remain loaded in accelerator memory, which makes failover fast at the expense of a device that serves nothing while the primary is healthy.

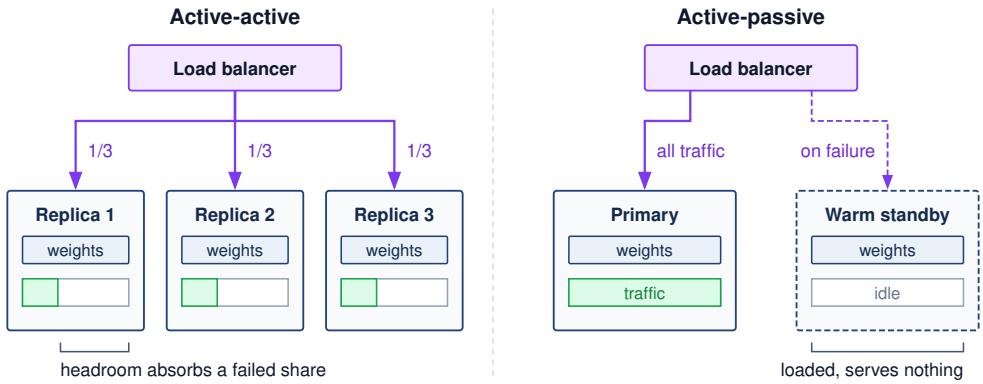

Figure 7.17: Serving Redundancy Strategies: Active-active (left) splits traffic across three replicas, each carrying about one-third, which uses every device but requires enough headroom for the survivors to absorb a failed replica's share. Active-passive (right) sends all traffic to one primary and keeps a warm standby with weights already loaded, which simplifies failover at the cost of capacity that serves nothing while the primary is healthy.

Geographic replication extends the same choice across regions. It protects against data-center or regional network failures, but requests routed to distant regions pay additional latency, and the active-active headroom problem reappears at regional scale. When a region fails, the surviving

regions must absorb its traffic and, for stateful services, rebuild its sessions, so each region either carries spare capacity sized for the largest region it may absorb or accepts degraded latency during the failover. Traffic returns to a recovered region gradually, because its caches and warm replicas start empty and an abrupt shift back would overload it and trigger the same failover again. Multi-tier designs combine several replication policies at once: edge caches replicate for latency, regional serving clusters replicate for availability, and a global primary or coordination layer preserves consistency and freshness.

7.11.2.3 Replica placement and failure domains

Effective redundancy requires placing replicas in independent failure domains. Different machines tolerate individual machine failures; different racks tolerate rack-level failures from power or top-of-rack switch issues; different availability zones, the separately powered and networked sections of one region, tolerate data-center-section failures; and different regions tolerate entire data-center failures. The independence level should match the availability requirement and cost constraint. Regional replication is expensive because it duplicates compute and network capacity, but it is necessary when regional failure must not become user-visible downtime. Placement alone, however, only creates the possibility of recovery; the serving system still needs a control loop that notices the failure, removes the affected replica from traffic, and preserves enough state that the user-visible service continues with bounded degradation.

7.11.3 Failover mechanisms

When a replica fails, traffic must be redirected to healthy replicas, and the speed and reliability of that redirection determine how much of the failure users see. The mechanism decomposes into two questions: how the system decides that a replica can no longer serve, and what happens to stateful sessions already attached to it.

7.11.3.1 Health checking

Health checks decide when traffic should move. The liveness and readiness probes of section 7.9.1 apply to a serving replica, where readiness means that its weights are loaded, its accelerator is initialized, warmup has completed, and dependencies such as feature stores (the services that supply precomputed model input features) are reachable. Serving adds an inference check, which runs a known input through the model and compares the output with the expected result, catching the silent failure in which the process is healthy but its answers are wrong. The load balancer acts on these signals.

The probe parameters set the failover budget. With a 5 s check interval, a 2 s timeout, and a failure threshold of 3 missed probes, a crashed replica leaves rotation 12–17 s after it fails, depending on where the crash falls between probes. Every request routed to it in that window fails or waits for a client retry. Replica-level failover on this time scale therefore protects subsequent traffic, not the requests that hit the crash; a request with a millisecond budget can still meet it only through a retry on another ready replica within its remaining deadline. Shortening the interval or lowering the threshold speeds failover but raises the false-positive rate, marking healthy replicas unhealthy during transient stalls and removing capacity exactly when load is high. A success threshold of 2 clean probes applies the same caution in reverse, holding a recovered replica out of rotation until it has proved stable.

↪ What each probe must test on an accelerator replica, from process liveness to a fixed-prompt inference check, is developed in section 10.11.9.

7.11.3.2 Session affinity and stateful failover

For stateful serving, session affinity routes all requests within a session to the same replica. Load balancers maintain session-to-replica mapping through sticky sessions implemented with cookies, headers, or IP hashing. Affinity preserves locality but supplies no replacement copy when that replica fails. Checkpointing saves session state periodically, replication copies it to a standby, and external systems such as Redis or Memcached move it outside the model worker. The external store still needs its own availability and durability contract. Accepting state loss is valid only when the service can reconstruct derived state from retained history or explicitly begin a reduced-context session. Which option fits a workload depends on state size, update frequency, and the quality impact of state loss, as table 7.13 summarizes across the four common approaches.

Table 7.13: Stateful Failover Approaches: Recovery latency, state guarantees, and operational complexity distinguish four approaches to session-state failover. Checkpoint recovery may lose updates after the saved snapshot. Distributed storage can provide different consistency contracts, and synchronous replication provides its promised guarantee only under the specified acknowledgment and fault model.

Approach	Recovery latency	State guarantee	Operational complexity
State loss	Fast	None	Low
Checkpointing	Medium	Consistent snapshot; later updates may be lost	Medium
Synchronous replication	Fast	Strong	High
Distributed state	Fast	Configurable	Medium

7.11.4 Model-specific serving fault tolerance

Retrying a failed image classification request on another replica loses no session state, whereas losing the replica that holds a multi-turn LLM conversation drops gigabytes of live attention state and losing a recommendation feature store leaves the ranker with stale user signals. Model-specific serving policies therefore start from a state-loss budget: which state can be rebuilt within the latency deadline, which state must be replicated in advance, and which state may be degraded temporarily without violating the product contract. Each of the three common serving regimes answers that budget differently, setting both its fault-tolerance policy and its degradation path (table 7.14).

Table 7.14: Serving Recovery State Invariants: Serving fault tolerance starts from the state whose loss changes user-visible quality. LLMs protect live context, recommendation systems protect feature freshness, and vision services protect preprocessing and device correctness.

Serving workload	State expensive to lose	Fault-tolerance policy	Degradation path
LLM conversation	KV cache, conversation context, prefix state	Replicate or checkpoint high-value session state; regenerate only when the latency budget allows.	Rebuild from transcript, reuse cached prefixes, or ask for retry.
Recommendation	Fresh user features, item features, embeddings	Replicate feature stores, cache hot features, and monitor freshness as a reliability signal.	Serve stale or default features with measured quality loss.
Vision service	Preprocessing state, device health, edge input	Retry stateless requests on healthy replicas and fail closed on preprocessing or accelerator-health faults.	Use a smaller local model, defer prediction, or return low confidence.

The KV cache[19] can be substantial (gigabytes for long contexts across attention layers). Losing it requires recomputing attention state from the retained token history, which can take seconds to minutes; previously delivered text need not be generated again. During an in-flight response, recovery must also know which output tokens the client has already received so a retry does not silently duplicate or replace part of the response.

LLM serving choices trade KV-cache recovery latency against memory and storage overhead. The simplest strategy is to accept regeneration cost by rebuilding KV cache from conversation history after failure, but this can turn a long conversation into a multi-second or multi-minute restart. KV-cache checkpointing saves live state periodically and bounds the amount of regeneration, while KV-cache replication keeps a standby copy for fast failover at the cost of extra memory. Prefix caching, also offered as prompt caching, reuses state for common system prompts and shared context. A compatible cached prefix can reduce reconstruction work if it survives the failure and remains accessible to the replacement worker. Cache reuse alone does not establish a failover guarantee or preserve the session-specific history.

Recommendation serving protects freshness rather than conversational continuity. Recommendations depend on user and item features from feature stores, so feature-store unavailability can either degrade ranking quality or block recommendations entirely. The serving decision is the staleness budget. Replicated feature stores keep recent features available across zones, local caches cover hot features, fallback to stale features accepts measured quality loss, and default features let the service return a lower-quality result when lookup fails. Real-time features such as recent user actions make freshness monitoring part of fault tolerance, because a pipeline that keeps serving old features is available operationally but wrong semantically. Large embedding tables may also sit behind dedicated embedding services, which need their own replication and failover plan.

Vision serving is usually closer to the stateless end of the spectrum because each image or frame can be retried on another replica. That simplicity changes where the checks belong. Thermal throttling threatens latency, while uncorrected memory errors threaten numerical validity; health policy must distinguish those consequences. Preprocessing checks must also prevent invalid resize, crop, color conversion, or normalization from producing apparently valid predictions. When a preprocessing or accelerator-health check fails, the service fails closed, withholding the prediction rather than returning one it cannot trust. At the edge, disconnection can require a validated smaller local model, delayed nonurgent work, or abstention when the available computation cannot meet the prediction contract.

[19] **KV cache**: Stores key and value projections for all previous tokens across all attention layers, with a per-token size derived in section 10.7.1. For a Llama-family 70B model at its 128K context in BF16, the KV cache is about 43 GB per conversation. Reconstruction can exceed serving latency service level agreements (SLAs) even when a replacement replica is already available.

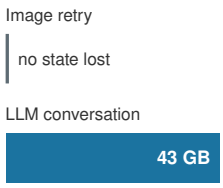

A retried image loses nothing; a lost conversation loses 43 GB of KV cache.

When redundancy cannot preserve full service, the remaining choice depends on the application contract. A validated reduced-capability response can preserve useful availability; if no such response is acceptable, the system must reject or defer the work. Graceful degradation makes the permitted reduction explicit before failure occurs.

7.12 Graceful Degradation

During a major regional network outage, an e-commerce site suddenly loses access to its heavy, GPU-accelerated recommendation cluster. Rather than showing users empty pages or crashing, the site instantly switches to serving precomputed, generic popular items. The serving fault tolerance mechanisms developed in section 7.11 aim to maintain full service, but graceful degradation dictates what happens when those defenses are overwhelmed.

Definition 7.3: Graceful degradation

Graceful degradation is a fault tolerance strategy in which a system responds to resource exhaustion or component failure by deliberately reducing service quality (falling back to a smaller model, serving cached results, or returning partial outputs) rather than failing completely, maintaining measurable availability at reduced capability.

1. **Significance**: Graceful degradation converts total outage risk into a controlled quality reduction. A recommendation system can fall back from a large ranker to a collaborative-filtering model on GPU failure, maintaining request availability with a quality penalty that must be measured for that workload. The engineering calculation includes both the requests served and the utility lost during fallback.
2. **Distinction**: Unlike complete system failure (where the service returns errors to all requests), graceful degradation provides a managed transition through a predefined capability hierarchy in which each fallback level has known accuracy, latency, and resource characteristics that allow the system to continue satisfying SLOs at reduced quality.
3. **Common pitfall**: A frequent misconception is that degradation is automatic in well-designed systems. Graceful degradation requires explicit preengineering. Fallback models must be preloaded or precomputed, switchover logic must detect the failure condition and trigger the transition without human intervention, and each degradation level must have been validated to satisfy its reduced SLO before it is needed in production.

7.12.1 Degradation dimensions

A degradation plan must name which service property will be sacrificed before an incident begins. The main levers in table 7.15 form a pre-incident design checklist, not an outage-time brainstorming aid. Each sacrificed property needs a measured quality budget, an activation condition, and a recovery path before the failure occurs.

Table 7.15: Graceful Degradation Dimensions: Degradation plans choose which service property can be reduced while still returning a useful answer. Naming the dimension in advance prevents the system from treating every failure as a binary serve-or-crash decision.

Dimension	What the system sacrifices	Typical fallback
Quality	Model accuracy or ranking quality	Use a simpler model, such as a collaborative-filtering fallback for a multi-tower recommender.
Latency	Response speed	Batch more requests together to preserve throughput and model quality under high load.
Coverage	Result completeness	Return top-10 search results instead of top-100.
Freshness	Recency of computed results	Serve cached or stale candidates, such as hour-old news recommendations during an outage.

Dimension	What the system sacrifices	Typical fallback
Feature completeness	Input richness	Use content features when user history or real-time context is unavailable.

These dimensions are not interchangeable. A search product may sacrifice coverage by returning fewer results while keeping freshness intact. A recommender may sacrifice freshness by serving cached candidates while preserving latency. An interactive assistant may sacrifice quality by using a smaller model while keeping the session alive. The correct fallback is therefore workload-specific. It follows the property whose temporary loss least harms the application and whose recovery path can be validated before the incident.

7.12.2 Graceful degradation strategies

The dimensions in table 7.15 become operational only when the service binds each sacrificed property to a specific fallback mechanism. The strategy ladder starts with model fallback when the primary inference path is too expensive or unavailable. It then moves to feature fallback when input quality degrades, and finally to load shedding when demand exceeds capacity and the system must preserve the highest-value requests first.

7.12.2.1 Model fallback

Model fallback turns quality into an explicit availability lever by maintaining multiple model versions with different resource requirements. The primary model provides full capability at the highest resource cost, a secondary model reduces capability and resource demand, a tertiary model preserves minimal capability, and a static fallback serves precomputed defaults with no inference at all. When the primary path is unavailable or overloaded, the policy descends only as far as resource pressure requires, then returns upward after the primary path is healthy again.

In an image classification cascade, the primary path might run a ViT-Large model with 307 million parameters and 88 percent ImageNet top-1 accuracy. Under pressure, it might fall back to EfficientNet-B4 with 19 million parameters and 83 percent accuracy, then fall back again to MobileNet-V3-Large with 5.4 million parameters and 75 percent accuracy when only minimal classification remains affordable. If even that path cannot meet the SLO, the service can return cached labels, a coarse category, or "classification unavailable" rather than blocking the request indefinitely. For that ladder to be real rather than aspirational, the production system must keep multiple models deployed, route requests to the appropriate level, monitor fallback frequency, and measure the quality impact of each level.

Recommendation and ranking systems use the same pattern when a complex deep learning model falls back to a top-n popularity list, a linear ranker, or a cached response during primary serving failures. Complex models often fight for the last mile of accuracy while simple heuristics provide the bulk of the utility, so the fallback path must be monitored as a first-class quality signal rather than treated as a silent failover. Otherwise the product appears healthy while quality is silently traded for availability.

7.12.2.2 Feature fallback

Feature fallback controls how much input quality the model may lose before the request must be blocked. When feature retrieval fails, validated population-level defaults or cached values can substitute for missing user or item features. A recommender might use genre-level item defaults for missing item attributes or the most recent cached value of a real-time signal. The model must have been trained or evaluated with those missing-feature semantics; an average embedding or a vector from another encoder is not automatically a compatible input. Approximate features are useful only when their representation and measured prediction quality satisfy the fallback contract.

The key design step is prioritizing features by their measured contribution to prediction quality. In a four-level policy, critical features block the request, important features use defaults, useful features use cached values, and optional features can be omitted (table 7.16). The table's quality-loss ranges are example policy budgets, not transferable properties of those feature categories. Deployment requires measurements showing that each permitted fallback meets its budget. When

feature degradation alone is insufficient to maintain system stability under extreme load, the serving infrastructure must begin actively dropping requests.

Table 7.16: Graceful Feature Degradation Tiers: Features classified by importance to prediction quality, with the corresponding response when each tier is unavailable. Critical features block the request entirely; lower tiers allow the system to fall back to defaults, cached values, or omission with bounded quality impact. The tiers convert a binary "feature available or not" question into a four-step degradation ladder.

Tier	Example features	Missing action	Quality impact
Critical	User ID, Item ID	Block request	Cannot serve
Important	User history, Item attributes	Use defaults	5–10% quality loss
Useful	Real-time context	Use cached	2–5% quality loss
Optional	Secondary signals	Omit	$< 2\%$ quality loss

7.12.2.3 Load shedding

Load shedding protects the system by sacrificing selected requests before overload spreads to every request. Random shedding is the simplest policy: drop a fraction of incoming requests and preserve enough capacity for the rest. Its weakness is that it treats all requests as equally valuable. Priority-based shedding turns quality-of-service degradation into an explicit policy by serving premium users before free users, revenue-generating requests before analytics requests, and interactive traffic before background batch work. The system is still degrading, but it is degrading according to a predeclared service contract rather than whichever queue happens to fill first.

Admission control applies the same idea at system entry points. It rejects requests that would exceed capacity rather than accepting work that will degrade every request already in the system. Circuit breakers extend the same protection to individual dependencies by failing fast when a model replica, feature store, or supporting service is unhealthy, so that one failing dependency cannot exhaust connection pools across the fleet. Section 10.11.10 develops their states and the backpressure signals that accompany them.

7.12.2.4 Graceful degradation implementation

Implementation begins by turning the degradation ladder into a control loop. Tail-latency percentiles reveal delayed responses, dependency-specific error rates locate failed calls, resource utilization exposes pressure, and growing queues indicate a mismatch between arrivals and completions. These signals support entry and exit conditions for each degradation level. A trigger can justify load shedding before diagnosis is complete, even though the signal alone does not establish the underlying cause. Degradation triggers define conditions that activate degradation; three common triggers watch tail latency, error rate, and feature store health, and each one activates its own fallback mechanism (listing 7.1).

Listing 7.1: Progressive Degradation Triggers: Conditions that activate graceful degradation based on tail latency, error rate, and feature store health. Each trigger activates a different fallback mechanism appropriate to the detected problem.

```
# Monitor tail latency for user-facing impact
if p99_latency > threshold:
    activate_model_fallback()  # Switch to faster, simpler model

# Track error rates to detect downstream failures
if error_rate > threshold:
    activate_circuit_breaker()  # Fail fast, prevent cascade

# Feature store slowdowns degrade recommendation quality
if feature_store_latency > threshold:
    activate_feature_fallback()  # Use cached/default features
```

Degradation should increase progressively as conditions worsen rather than switch at one cliff. The system can raise the fallback percentage gradually as load increases, then require sustained im-

provement before traffic returns to higher-capability paths. Hysteresis prevents oscillation between primary and fallback modes when the service hovers near a threshold.

7.12.3 Degradation monitoring and alerting

Because graceful degradation trades quality for availability, monitoring must make that trade visible rather than letting the system appear healthy while serving lower-fidelity results. Table 7.17 turns that requirement into an action map: fallback share reveals how often model quality was reduced, default-feature share reveals input-quality loss, drop rate reveals capacity protection, and the primary-vs-fallback quality gap reveals whether the degraded path remains acceptable. Its escalation thresholds illustrate one policy; the service's permitted degradation duration and quality loss must determine the deployed thresholds.

Table 7.17: Degradation Monitoring Action Map: Graceful degradation is safe only when the quality trade is observable. Each signal ties a degraded serving mode to the escalation and postincident question that prevent silent quality loss.

Signal	Quality trade exposed	Escalation rule	Postincident question
Fallback-model request share	Model capability reduced to preserve serving	Sustained activation beyond 5 minutes becomes a warning.	Did the fallback preserve enough user value?
Default-feature request share	Input richness reduced because dependencies lag	Severe degradation affecting over 50% of requests becomes critical.	Which feature dependency needs replication or caching?
Load-shed request rate	Coverage reduced to protect the service	Any growth after the circuit breaker opens triggers capacity investigation.	Was admission control early enough to prevent cascade?
Primary-vs-fallback quality gap	Product quality lost during the incident	Degradation beyond 1 hour escalates even when availability metrics look healthy.	Is the fallback path still a valid product experience?

After the event, postincident analysis should identify the root cause, evaluate fallback effectiveness, measure user and business impact, and define improvements that prevent recurrence. Implementing these fallback mechanisms safely requires deep visibility into the system's runtime state. When a complex recommendation pipeline begins degrading, operators must rapidly determine which microservice is failing, and that diagnosis depends on distributed debugging and observability.

7.13 Distributed Debugging and Observability

At 2:00 AM, latency on the flagship generative API spikes from 200 ms to 5 seconds, but CPU, memory, and network metrics all look normal. The latency breach can justify immediate mitigation under a predefined policy; choosing a targeted repair requires locating the stalled work across hundreds of microservices. Distributed debugging and observability provide the diagnostic capability required to locate these invisible bottlenecks.

7.13.1 Diagnostic challenges in distributed ML systems

When a latency spike or stalled collective appears on one run and vanishes on the next, the evidence needed for recovery is split across timing, state, scale, and model behavior. Nondeterminism enters from multiple directions: network timing variations change execution order, thread scheduling differences alter race conditions, GPU kernel execution order varies across runs, and floating-point operation ordering changes numerical results. Consequently, a bug that manifests on one execution may not reproduce on subsequent runs, creating "Heisenbugs"[20] that appear to vanish when observed.

Partial failures turn that nondeterminism into a recovery problem. A single machine can also suffer partial failures, but distribution adds independently progressing workers and communication paths. One component can stop while others continue, leaving peers uncertain whether it has failed or become unreachable. Scale then removes manual inspection as a viable debugging method. With thousands of components, automated tools must filter large telemetry streams and identify the few signals that explain the incident. ML systems add model behavior to the same diagnostic problem: silent accuracy degradation produces wrong results without errors, numerical issues like NaN and infinity propagate through computation, data-dependent bugs manifest only for specific inputs, and

[20] **Heisenbug**: A pun on Heisenberg's uncertainty principle, popularized in Gray (1986)'s analysis of why production systems fail. Instrumentation perturbs timing, so observing the bug changes it. In distributed ML training, inserting gradient logging alters collective timing and can mask the very race that caused the silent corruption.

learning causes expected behavior changes that can resemble bugs. Isolating a fault across all four sources requires structured telemetry that links cluster-wide symptoms to individual requests and events.

7.13.2 Observability pillars

Observability supports both immediate mitigation and investigation of the cause. Suppose a recommendation service suddenly violates its latency SLO while the model replicas still report healthy GPU utilization. Metrics establish the timing and scope of the symptom, traces show which service span consumes the request budget, and logs record events inside that service. Figure 7.18 connects these streams to operational decisions. A single signal can justify an action, such as shedding load when a queue exceeds its budget. Joining request and trace context with model versions, deployment tags, and time windows strengthens the diagnosis when several causes could explain the same symptom. Correlated evidence narrows those possibilities; it does not by itself prove the physical or software cause.

↪ How metrics, logs, and traces become a fleet-wide observability architecture that routes evidence across many models and owners is developed in section 12.6.6.

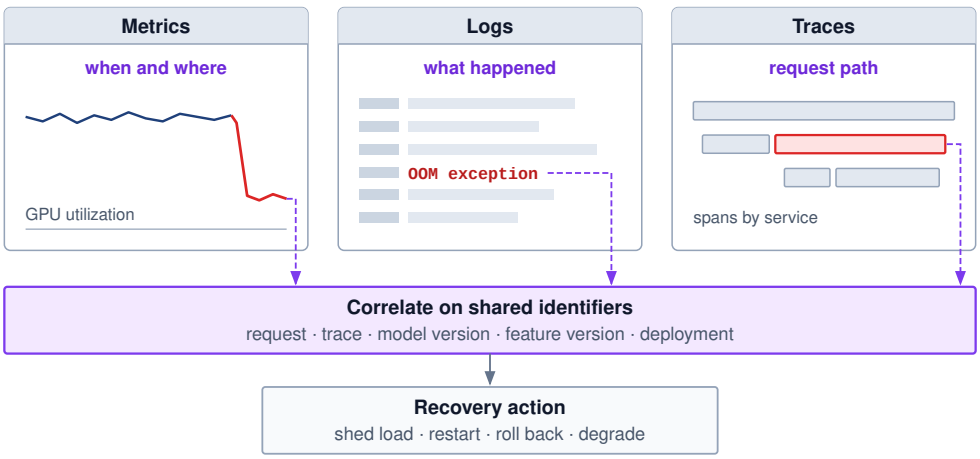

Figure 7.18: The Three Pillars of Observability: Diagnosis of fleet failures can combine three signal types. Metrics reveal *when* and *where* a problem occurs (for example, GPU utilization drops). Logs provide the *what* through event context (for example, an out-of-memory exception). Traces expose the *request path* by linking events across the distributed lifecycle (for example, identifying the request that triggered the memory spike). A single signal can justify immediate mitigation; joining signals on shared identifiers strengthens diagnosis without proving the cause.

Metrics are the first evidence because they compress fleet behavior into time series that can trigger action. A metric stream becomes fault-tolerance infrastructure only when an alert maps to a recovery decision. High p99 latency may trigger load shedding, a falling cache hit rate may trigger fallback features, and a sudden prediction-distribution shift may block a rollout before it corrupts user-facing decisions.

Logs provide the event context that a metric intentionally discards. Structured logging uses formats such as JavaScript Object Notation (JSON) with consistent fields, so incident tooling can search by service, request, model version, GPU, feature store, or trace identifier. A GPU memory allocation failure logged this way records both the local failure context and the trace identifiers needed to connect the event to distributed metrics and spans (listing 7.2).

Listing 7.2: Structured JSON Log Entry: A structured log entry for a GPU memory allocation failure, including trace and span identifiers for correlation with distributed traces and resource metrics for diagnosis.

```
{
  "timestamp": "2024-01-15T10:23:45.123Z",
  "level": "ERROR",
  "service": "inference-server",
  "trace_id": "abc123",
  "span_id": "def456",
  "message": "GPU memory allocation failed",
  "gpu_id": 3,
  "requested_bytes": 4294967296,
  "available_bytes": 2147483648
}
```

For fault tolerance, the important logging property is consistent correlation rather than the particular backend. Log levels should map incidents to escalation consistently across components, but the recovery system depends more on whether operators can query events by the same request, trace, model version, feature version, and deployment identifiers that appear in metrics and spans. Log aggregation is therefore useful because it preserves the join between local failure context and fleet-wide symptoms.

Traces show where latency or failure propagated across distributed components. A trace is the end-to-end journey of a request, a span is one operation within that journey, and context propagation carries the trace identity across service boundaries. A trace through an ML inference pipeline illustrates how these spans compose:

```
Trace: user-request-12345
|-- Span: api-gateway (5 ms)
|   `-- Span: auth-service (2 ms)
|-- Span: feature-service (15 ms)
|   |-- Span: user-feature-lookup (8 ms)
|   `-- Span: item-feature-lookup (12 ms)
|-- Span: inference-service (45 ms)
|   |-- Span: preprocessing (3 ms)
|   |-- Span: model-inference (40 ms)
|   `-- Span: postprocessing (2 ms)
`-- Span: response-formatting (1 ms)
Total: 66 ms
```

Standard distributed tracing APIs such as OpenTelemetry carry this context across services, and a tracing backend stores and visualizes the traces. The durable requirement is not the particular backend but enough preserved context to answer whether the fault lives in request routing, feature retrieval, model execution, postprocessing, or a dependency outside the model service.

A single training job uses the same three pillars with the job, step, and rank as its shared identifiers. Its metrics are per-rank series such as step time, gradient norm, and loss, so a rank that lags or diverges stands out against its peers. Its logs are each rank's records of errors, restarts, and collective timeouts. Its trace of a training step records when each rank entered and left each collective, which shows the rank the others waited on; Meta used such communication traces to locate stalled collectives and slow workers during Llama 3 pretraining (war story 7.2). Joined on those identifiers, the three streams let operators match a symptom to its first recovery action in table 7.7.

7.13.3 ML-specific debugging

Operational signals alone are not enough when a process stays healthy while its tensors, inputs, or per-rank execution times drift outside the model's contract. ML systems therefore add specialized debugging across three surfaces: numerical health inside the computation, semantic validity across the data pipeline, and straggler attribution across distributed stages.

7.13.3.1 Numerical debugging

Numerical debugging targets failures where tensors remain valid objects but contain invalid values. NaN detection[21] exposes values that can propagate through later arithmetic without a process crash.

[21] **NaN propagation**: Many IEEE 754 arithmetic operations propagate a NaN operand, so an invalid gradient can contaminate dependent updates without stopping the process. The affected scope depends on the computation's data dependencies. Zero divided by zero and the logarithm of a negative real value produce NaNs; division of a nonzero finite value by zero or floating-point overflow can instead produce infinity, which may cause a later invalid operation. Checking only for a process crash therefore misses numerical failures.

A minimal diagnostic check only flags the invalid output (listing 7.3), so the caller must act on the result and block its publication.

Listing 7.3: NaN Detection: This check flags NaN output for diagnosis. The caller must block publication or choose a fallback; a logged request identifier refers to an input retained by the application for replay.

```
# Check tensor for NaN values
# (propagate from any corrupted computation)
if torch.isnan(output).any():
    log.error("NaN detected in output", request_id=request_id)
    # Caller must retain a replayable input reference and block or replace this output
```

During training, per-step gradient statistics expose numerical faults before they corrupt the entire run. Tracking the global gradient norm catches sudden explosion or vanishing, histogramming the gradient distribution flags outlier spikes, and logging layer-wise gradient norms isolates which block in the network first destabilized. Mixed-precision training[22] can introduce related numerical issues when narrow dynamic range interacts with unstable updates. A dynamic loss scaler reduces its scale after detecting nonfinite gradients, indicating overflow or another invalid operation rather than proving underflow. Gradient magnitudes, skipped updates, and comparison with a higher-precision replay help distinguish range problems from other causes; ordinary FP16 rounding differences are not themselves faults.

[22] **Mixed-precision numerical issues**: FP16's exponent range is far narrower than FP32's or BF16's, which is why mixed-precision training scales the loss before backpropagation and dynamic scaling skips nonfinite updates. The formats and the mixed-precision recipe are recapped in section A.6.

7.13.3.2 Data debugging and validation

Data bugs surface as valid-looking but wrong inputs, so debugging must verify semantic correctness at each pipeline boundary in addition to numerical validity. Validation checks expected format, shape, value ranges, required fields, and encoding before data reaches the model, catching malformed examples while the failure is still local to the input path. Once the input passes those static checks, the system needs evidence about behavior over time. Distribution monitors track feature drift, null-rate changes, and outliers as they emerge, while transformation instrumentation logs intermediate shapes and statistics so teams can compare each stage against known-good data. Together, these signals separate a corrupted input path from a numerical failure inside the model.

↪ The validation gates and response policies that run these data checks continuously across a production fleet are developed in section 12.4.7.

7.13.3.3 Straggler detection and analysis

Perspective 7.6 treated a straggler as a slow training rank that holds up a synchronous step. The same idea extends to serving once "peer" means any component on a request's path, so straggler diagnosis there asks which replica or service stage is slow enough to become a failure-equivalent bottleneck. A slow serving replica delays requests routed to it and can trigger retries under load; its effect on the service's tail percentile depends on its traffic share and the routing policy. A request that waits for several required parallel branches is instead bounded by the last required completion. Operation-level timing instrumentation measures time for each operation, enabling latency attribution across pipeline stages. Wrapping each operation in a context manager that emits per-span timing metrics makes performance directly comparable across replicas (listing 7.4).

Listing 7.4: Operation-Level Timing: Context manager instrumentation attributes latency to individual pipeline stages, enabling straggler detection by comparing the same span across replicas.

```
# Wrap operations in timing context managers for latency attribution
with timer("feature_lookup"):
    features = feature_store.lookup(
        ids
    )  # Often the latency bottleneck
with timer("model_inference"):
    predictions = model(features)  # Host submission time for asynchronous GPU work
```

For asynchronous GPU execution, the host span around `model(features)` in listing 7.4 measures only the CPU time required to enqueue kernels onto the GPU stream, not completed device work.

Because kernel launches return to the host before the GPU finishes executing them, a naive host timer can report a fast inference span even on a thermally throttled GPU and misattribute the delay to whichever later stage first synchronizes the stream or copies predictions back to host memory. Measuring actual GPU execution time requires recording hardware events on the stream before and after the forward pass and reading their elapsed time after the stream completes. Like-for-like component timing is then compared across replicas with percentile analysis. The p50 latency shows typical performance, while p99 shows tail latency, and comparing p99 across replicas identifies workers whose slow path has become a fleet-wide bottleneck. The diagnosis must then separate the causes that share this slow-worker symptom, which include the gray-failure causes of section 7.9.5 and data skew, in which some inputs are slower to process.

7.13.4 Serving failure patterns

Serving failures require a symptom-to-action mapping under much tighter latency constraints than the training-loss signatures of section 7.9.1. Common serving symptoms differ both in what each one reveals and in what the serving system must do first (table 7.18). Cascade failures are hardest to diagnose because the root cause is obscured by the secondary timeouts and retries it generates, as example 7.1 traces across metrics, logs, and spans.

Table 7.18: Serving Failure Response Map: Serving incidents must be classified quickly because every diagnosis competes with a user-facing latency budget. The table ties each symptom to the first operational response rather than treating all serving failures as generic errors.

Symptom	Likely interpretation	Immediate response
Latency spikes	Resource contention, garbage collection, cold caches, or model reloads	Check placement and capacity; repeated spikes usually mean the problem is no longer transient.
Rising error rates	Dependency failure, data-format change, or model bug	Investigate immediately because errors compound across dependent services.
Silent quality degradation	Model drift, feature degradation, or data-pipeline fault	Use quality monitoring beyond standard operational metrics.
Cascade failure	Timeout exhaustion, resource depletion, or error propagation from one failing component	Trip circuit breakers and preserve isolation boundaries before the failure spreads.

Example 7.1: Diagnosing a cascade failure

Scenario: A recommendation system experiences sudden latency spikes.

Diagnosis: The three observability pillars narrow the failure path. Metrics reveal the symptom: p99 latency jumped from 45 ms to 800 ms at 14:32, with error rate increasing from 0.1 percent to 15 percent. Traces locate the delay in the feature-service span, which consumes 700 ms instead of the normal 12 ms, while model inference remains at 40 ms. Logs show repeated connection timeouts to the user embedding cache at 14:31, followed by requests bypassing the cache and reaching the embedding database directly. That evidence identifies cache unavailability and increased downstream demand; it does not yet distinguish a failed cache process from a network or resource problem.

Systems lesson: Cache unavailability can transfer enough demand to overwhelm the embedding database. A circuit breaker and a validated default-embedding path can limit that cascade while diagnosis continues. The fallback must also respect the quality budget and avoid sending the rejected demand to another saturated dependency.

Observability and diagnostic maps help operators isolate a failure once it strikes, but verifying that detectors, circuit breakers, and recovery paths actually hold under fire requires triggering those failures deliberately before production traffic meets them.

7.14 Fault Injection Tools and Frameworks

An injected failure is useful evidence only when it resembles the production failure the system must survive. A multi-node training cluster that claims to tolerate network partitions should have that connection severed in a staging environment so engineers can observe what the orchestration layer does. Fault injection is the engineering discipline of deliberately perturbing a system to empirically verify that robustness mechanisms work before production traffic discovers the gap. Fault-injection tools for tensors and accelerators cover bit flips, corrupted tensor values, and accelerator-oriented failures (Z. Chen et al. 2020; Tsai et al. 2021; Gräfe et al. 2023), while network partitions, latency injection, dependency termination, and API-response faults belong to chaos and distributed-systems testing practice (Basiri et al. 2016).

An injection experiment is evidence only at the level where it injects. Its fault model names the originating disturbance, such as a transient single-bit upset or a permanent multi-bit defect, and its error model names the corrupted state that results. Results for transient single-bit faults therefore do not establish resilience to permanent multi-bit defects (Sangchoolie et al. 2017). The injection level also decides which masking the experiment includes. Hardware absorbs many physical disturbances before software sees them, in unused units, overwritten values, and dead instructions (Mohanram and Touba 2003; Ko 2021). Hardware-based injection through field-programmable gate array (FPGA) emulation or particle-beam exposure (Velazco et al. 2010) passes through that masking, while software tools such as PyTorchFI (Mahmoud et al. 2020) perturb weights, activations, and gradients directly. They cover many model states quickly, but by skipping the masking they can overstate how often a physical fault corrupts the model. Frameworks such as Fidelity (He et al. 2020) narrow that abstraction gap by mapping hardware faults to the software-visible errors they produce. A tensor mutation still tests how the software responds to a corrupted value; on its own, it does not estimate production vulnerability.

For fleet recovery, the most valuable experiments inject faults at the level where recovery decisions are made. Killing a rank in the middle of a collective or delaying its heartbeat past the timeout exercises the detection layers of section 7.9.1. Partitioning a node from the rendezvous service exercises the membership and fencing rules of section 7.9.2. Corrupting one checkpoint shard exercises the manifest check of section 7.8.5.2. Each of these experiments tests one row of table 7.4. Running these experiments continuously in staging, and in controlled form in production, keeps rarely executed recovery code from decaying between real incidents (Basiri et al. 2016). A recovery path is trustworthy only for the fault model it has been exercised against, which is why the case studies that follow report both the failures production fleets met and the machinery that absorbed them.

7.15 Case Studies

Training a 175-billion-parameter language model across 992 GPUs for two months, keeping a 4,096-chip supercomputer slice routable despite daily hardware faults, and serving low-latency recommendations across dozens of microservices each stress a different layer of the fault-tolerance stack. Public operational reports from Meta, Google, and Netflix show how production fleets combine failure detection, state preservation, topology reconfiguration, and chaos testing at scale. Their reported numbers are operational magnitudes that expose the design pressure on each system rather than universal hardware constants.

7.15.1 Large-scale LLM training at Meta

Over two months of training the Open Pre-trained Transformer (OPT-175B) on 992 NVIDIA A100 GPUs, Meta's engineering team reported at least 35 manual restarts and estimated 70 or more automatic restarts triggered by hardware failures including ECC memory errors, NCCL and InfiniBand network faults, and lost GPUs (Zhang et al. 2022). Alongside these restarts, more than 100 hosts were cycled (Zhang et al. 2022; Meta AI Research 2022). These counts establish repeated interruption, but cycling hosts and restarting jobs are different events; their totals do not directly measure one common MTBF. Numerical instability created a separate recovery problem. The recovery design had to preserve progress while distinguishing failed execution resources from suspect optimization state.

During the OPT run, the team monitored progress through train-log freshness checks. The public logbook records a 15-minute modified-file threshold, later relaxed to one hour, as part of restart monitoring (Meta AI Research 2022). For a hang that raised no error, this threshold contributed to T_{detect}; polling and response delays could extend the wait before restart. The logbook illustrates the trade-off of equation 7.11 in operational form, since relaxing the threshold reduced sensitivity to delayed logs while allowing longer idle windows.

Recovery procedures also had to account for loss spikes associated with numerical instability. Unlike a fail-stop event, this symptom made recent model state suspect without identifying a failed device. The team combined rollback to an earlier checkpoint with a lower learning rate to recover from divergence (Meta AI Research 2022). The paper separately reports approximately 147 TFLOP/s per A100 (about 47 percent of dense FP16 peak) (Zhang et al. 2022); that execution-rate metric alone does not establish end-to-end useful throughput after interruptions and replay. The enduring lesson is the distinction between recovery paths. A failed device calls for replacement and restoration of valid state, while numerical divergence can require an older checkpoint and a changed optimization policy. Retaining checkpoint history makes that second path possible.

7.15.2 Google TPU pod resilience

Google's TPUv4 supercomputer combines configurable topology with recovery software. A pod contains 4,096 chips organized as 64 cubes of 64 chips, each arranged in a $4\times4\times4$ 3D mesh (Zu et al. 2024). Losing one chip removes $1/4096$th of pod capacity, but can interrupt an entire job that requires that participant. The recovery unit therefore exceeds the failed device.

War story 7.1 traced the recovery mechanism, in which the OCS fabric assembles healthy cubes into job-sized slices, reallocating spare cubes after machine outages and routing around unavailable links for supported OCS outage patterns. The published study adds the software that runs that loop, components including `libtpunet` and `healthd` for topology setup and health monitoring. It also gives the operating statistics that show how often the loop runs. System availability was 99.98 percent, with hardware outages affecting approximately 1 percent of training jobs (Zu et al. 2024). Daily failure rates were 0.08 percent of TPU machines, 0.005 percent of ICI cables, and 0.04 percent of OCS units. Small daily fractions still add up across thousands of components, so the control plane absorbs a steady stream of reconfiguration events rather than an occasional emergency. The architectural lesson is to match each failure boundary to a feasible replacement or routing mechanism.

7.15.3 Netflix chaos engineering for serving dependencies

While training resilience focuses on long-running batch jobs, Netflix's published work on Chaos Engineering (Basiri et al. 2016) provides the serving-side contrast, and it is the production origin of the injection discipline that section 7.14 applied to training. The premise is simple but operationally demanding. Because production systems have many interacting dependencies, teams should validate steady-state behavior by deliberately injecting failures instead of waiting for real incidents to discover whether fallbacks work.

For ML serving, the same discipline applies to feature stores, model-serving dependencies, caches, and fallback paths. A model endpoint can be healthy while a feature service is slow, stale, or unavailable; if the system lacks a tested fallback, a local dependency problem can become user-visible latency or degraded recommendations. Chaos tests make those assumptions executable by injecting latency, terminating dependencies, or forcing fallback paths in controlled conditions.

The recovery procedure for these injected faults centers on fallback hierarchies. If the primary deep learning model fails or times out, the serving system can switch to a lighter model, cached response, or simple popularity-based default. The systems lesson is that resilience in ML serving is not a static property but a continuous practice of active verification. A fallback path that is never exercised is documentation, not fault tolerance.

Across these production case studies, three architectural patterns recur. First, detection speed belongs in the recovery budget. Meta's log-freshness threshold and Google's health monitors determine when recovery can begin, while Netflix's chaos experiments test detection and fallback for selected dependency failures. Checkpoint age and replication lag separately govern recoverable state. Second, the atomic unit of failure shapes the architecture. Google combines cube replacement

with fault-tolerant routing, the OPT run restarted from durable checkpoints and kept their history for numerical rollback, and Netflix validates service-dependency fallbacks. Third, fault tolerance is a spectrum rather than a switch. Exact state recovery trades normal-operation overhead against lost work and recovery time; graceful degradation additionally trades service capability against continued availability. Those distinct trade-offs help identify the misconceptions that follow.

7.16 Fallacies and Pitfalls

An infrastructure team might spend weeks hardening their storage layer against disk failures, only to have their entire training run destroyed by a subtle software bug in their PyTorch distributed backend. Fault tolerance for distributed ML systems involves counterintuitive mathematics and subtle trade-offs where conventional data center wisdom often fails. Recognizing the fallacies and pitfalls below prevents engineers from optimizing for the wrong failure modes.

Fallacy: *Hardware redundancy is enough to keep a training run safe.*

This intuition comes from traditional systems, where disk failures, power outages, and network partitions dominate the incident mix. Hardware did cause most of the interruptions that the Llama 3 run recorded (figure 7.2), yet long-lived, stateful ML jobs also expose failure paths that hardware redundancy cannot cover. A malformed checkpoint path can make recovery impossible, a mismatched collective library can hang all ranks, an undersized shared-memory limit can crash data loaders, and a stale feature schema can silently corrupt training. Their relative frequency depends on the deployment, but even an infrequent software or configuration fault can destroy the same multi-day job that hardware redundancy was meant to protect. Investing heavily in hardware redundancy while neglecting software robustness (input validation, gradual rollouts, configuration management) leaves many important failure modes unaddressed, which is why the most reliable ML systems treat software bugs as inevitable and design defensively.

Pitfall: *Setting checkpoint interval by intuition.*

Organizations commonly set checkpoint intervals based on "feels right": "every hour seems reasonable" or "every 1,000 steps." The Young-Daly formula reveals these intuitions are often wrong. For the 1,000-GPU cluster of table 7.1, whose GPU-only MTBF is 50 hours, with a checkpoint write time of 5 minutes:

$$T^*_{\text{ckpt}} = \sqrt{2 \times 5 \times 3000} = \sqrt{30000} \approx 173 \text{ minutes}$$

The intuitive "every hour" checkpoints nearly three times more often than the expected-waste optimum. It reduces potential lost work, but the added write overhead outweighs that benefit under this model. A tighter RPO can still justify the shorter interval. The same hourly policy was too infrequent in the Young-Daly notebook, and both verdicts follow from one rule. The optimum scales as $\sqrt{T_{\text{write}} \cdot \text{MTBF}_{\text{system}}}$, so an hour is too long when a 24.5 s write meets a full-system MTBF of 3.87 hours, yet too short when a 5-minute write meets a GPU-only MTBF of 50 hours. If checkpoint time increases to 15 minutes (larger model, slower storage), the optimal interval stretches to 300 minutes, further still from the "every 15 minutes" that some teams adopt to "stay safe." The decision must balance expected waste with the recovery contract rather than treat either interval as universally correct.

Fallacy: *Component MTTF values predict individual failure timing.*

Engineers read a GPU MTTF of 50,000 hours and plan around a five-year replacement cadence per device. MTTF is a population mean, not a lifetime assigned to each component. A GPU with MTTF 50,000 hours might fail at 200 hours or last 100,000 hours. Under the chapter's independent constant-rate model, a fleet of 10,000 GPUs has an aggregate rate of approximately $10{,}000/50{,}000 = 0.2$ failures per hour, or roughly one failure every 5 hours. This planning model helps size recovery capacity and spare pools. It does not make device-health monitoring pointless. Temperature, error history, and recurring anomalies provide condition-specific evidence that the population mean alone lacks.

Pitfall: *Using MTBF calculations as if failures were independent.*

The reliability equation $\text{MTBF}_{\text{system}} = \text{MTTF}_{\text{component}}/N$ assumes component failures are statistically independent, while production failures correlate through the shared power, switches, software, and cooling of section 7.3.4.

Correlated failures change two quantities that the simple MTBF equation hides: incident frequency and blast radius. A cluster with 1,000 independent GPUs each with 50,000-hour MTTF has a first-GPU-failure MTBF of 50 hours. If a shared power, cooling, or software hazard adds an incident rate 10 times larger than that independent first-failure rate, the effective incident MTBF drops to about 5 hours. Separately, a correlated incident may take out 10 GPUs at once, increasing recovery cost even if the independent GPU failure rate itself has not changed. Reliability engineering must identify and mitigate correlation through diversity: different power feeds, different network paths, different software versions in canary deployments.

Fallacy: *Every failure should be handled as a full restart.*

Engineers often treat all failures as "node crashed, restart from checkpoint," but different failure modes require different responses. Transient stalls such as network congestion can permit a bounded wait or safe retry when state and the communication group remain valid; a brief fault can still require recovery if it corrupts state or leaves the group unusable. Permanent device loss requires replacement or a feasible resize, together with restoration of the missing state. Silent corruption requires a trusted recovery point, which may precede the latest checkpoint and therefore motivates checkpoint history retention. Persistent resource exhaustion (OOM, memory fragmentation) needs reconfiguration before restart; otherwise the job can crash again immediately. As section 7.3 details, misdiagnosing failure type wastes compute. Restarting for a harmless stall discards useful work, while continuing from corrupted state can contaminate later updates.

Pitfall: *Assuming checkpoints are consistent without validating state.*

Modern frameworks checkpoint transparently, creating the illusion of automatic consistency. In practice, distributed checkpoints require cross-node coordination that can fail subtly across multiple dimensions. Rank desynchronization occurs when workers record different training steps, such as rank 0 snapshotting iteration 1000 while rank 1 captures iteration 1001. Similarly, transient storage failures can cause partial writes that leave behind truncated shards, sharded optimizer state lag leaves optimizer states misaligned with model weights when captured asynchronously, and in-flight gradients from incomplete AllReduce operations can produce ambiguous reduction boundaries. Production systems must therefore implement checkpoint validation: verify all shards exist, verify iteration numbers match, and verify optimizer state matches model state. Without that validation, organizations that discover corrupted checkpoints during recovery from a failure have no recourse except restarting from a potentially much earlier checkpoint.

Fallacy: *Overhead budgets are fixed fractions of training time.*

Capacity plans often carry line items such as "pipeline bubble: 10 percent, checkpoint overhead: 5 percent, failure recovery: 3 percent," and such figures are easy to read as physical laws. They are engineering targets, not constants. Pipeline bubble overhead depends on the number of microbatches and the interleaved schedule, as its bubble fraction shows (section 6.5.2). Elastic recovery can reduce waiting for replacement capacity while still restarting workers and restoring state. Checkpoint overhead is a function of model size, storage bandwidth, and whether checkpointing is synchronous or asynchronous. Treating these numbers as fixed percentages leads to passive acceptance of avoidable inefficiency, because each one is set by a design choice.

Pitfall: *Using elastic training as a substitute for checkpointing.*

Elastic training adjusts parallelism degree when workers fail, but resizing still requires a consistent recovery state. Reducing the active worker count requires redistributing model shards, optimizer states, and data assignments, often through checkpoint reload. Below some minimum viable size, training becomes infeasible because the model no longer fits or the batch policy cannot be maintained; the job must wait for sufficient capacity before restoring. If a failure was caused by a software bug triggered by specific data, the bug persists in remaining workers. Elastic training is complementary to checkpointing, not a replacement. Resizing alone does not justify a lower checkpoint frequency, because the recovery contract, failure rate, and checkpoint cost still determine that choice.

Fallacy: *Silent data corruption is negligible because hardware has ECC.*

GPUs and memory systems include extensive error correction (ECC, CRC, parity), creating the intuition that silent data corruption is negligible. Large-scale CPU SDC and memory-error studies show that silent corruptions and memory faults still occur at fleet scale and can escape ordinary error reporting (Dixit et al. 2021; Sridharan et al. 2015). ECC and link CRCs remain necessary layers, but they protect stored and transmitted bits rather than arithmetic, so they are not an end-to-end correctness proof for a training run. A 10,000-GPU cluster executes enough arithmetic that even rare per-device corruption becomes an operating condition, which is why it needs the check-and-verify defenses of section 7.5 in addition to hardware protection.

Pitfall: *Adding GPUs without accounting for communication and recovery overhead.*

Capacity plans typically account for Amdahl's Law and AllReduce overhead but treat reliability as a fixed background condition. At extreme scale, each additional GPU increases the aggregate failure rate, which inflates the expected recovery and replay time per job. There exists a cluster size beyond which the time lost to failures and recovery exceeds the time saved by additional parallelism, and wall-clock training time increases rather than decreasing with N_{GPU}. This is the reliability version of diminishing returns, and it applies on top of the Amdahl ceiling, so the two effects must be modeled together. Whole-run useful goodput must count only retained progress and include recovery time in the denominator, and its reliability penalty generally grows with N_{GPU} under the independent-failure model. Section 9.2.1 composes that runtime loss with the scheduling and program losses without counting any of them twice.

7.17 Summary

Fault tolerance turns the statistical expectation of hardware failure into a manageable operational routine. Individual component reliability compounds across thousands of devices, driving system-level MTBF from years down to hours. At fleet scale, systems must bound lost work and recovery delay through recovery that is automatic, fast, and minimally disruptive to training or serving.

In training, progress is only as safe as the last state the fleet can prove it can restore. Checkpointing bounds lost work, and the Young-Daly interval (equation 7.7) prices that bound against write cost, shrinking only with the square root of a falling MTBF. Even at that optimum, under independent failures and a fixed write time, the share of compute lost to saves and rework grows with the square root of cluster size. RPO and RTO turn the expected-value optimum into a service contract. A manifest published last defines which generation exists, and fencing at the resource boundary decides who may still write to it. The retries that recovery triggers stay safe only when each operation is idempotent, inherits its caller's deadline, and draws on one shared retry budget, so the response cannot amplify the incident. Elastic recovery then trades capacity for continuity, valid only while the data, batch-size, and optimizer invariants survive the resize.

Serving inverts the budget. Training can spend minutes restoring state that satisfies its optimization contract; an interactive request has milliseconds, so serving buys availability with ready replicas and decides in advance which state it will rebuild, replicate, or abandon. The KV cache is the clearest case, derived state that can be reconstructed from retained history only when the latency budget allows. When neither redundancy nor reconstruction fits that budget, graceful degradation keeps the service useful by sacrificing a named property that was validated before the incident.

Recovery depends on evidence about both the failure and the state being restored. A loss spike can originate in hardware, data, numerical behavior, or synchronization. Suspect state can justify a pause

or rollback while investigation continues; attributing the cause to a device requires corroborating telemetry, replay, or integrity evidence. Restore drills and fault injection then test whether the selected response actually preserves valid state and meets its recovery budget. Passing those tests establishes evidence for the exercised fault model, rather than a guarantee against every failure.

Key Takeaways: Failure is normal operation

- **Scale makes failure routine**: A 10,000-GPU cluster encounters hardware failures every few hours under the chapter's component-rate assumptions; software must treat failure as a normal state.
- **Liveness is not correctness**: A gray failure keeps answering heartbeats while every peer waits on it, and silent data corruption passes ECC while feeding wrong gradients into a correct AllReduce, so detection needs progress measurements and check-and-verify evidence beside liveness checks.
- **Recovery is a contract**: RPO bounds lost progress, RTO bounds recovery delay, MTTR combines with MTBF to determine steady-state availability, a manifest identifies committed state, and fencing prevents an expired worker epoch from writing after replacement.
- **Checkpoint cadence balances two costs**: Under the first-order assumptions, $T^*_{\text{ckpt}} = \sqrt{2 \cdot T_{\text{write}} \cdot \text{MTBF}_{\text{system}}}$ balances checkpoint write overhead against expected rework; for the 175B running example on a 10,000-GPU cluster it lands near 13.8 min. The selected interval must also meet limits on lost progress; recovery delay needs its own budget.
- **Elasticity enables persistence**: Designing training jobs to be "elastic" (rebuilding communication groups and rebalancing data shards as traced in figure 7.16) can reduce waiting for replacement capacity. The runtime must preserve valid state and data coverage, and either preserve batch size or validate the changed optimization policy. Checkpoint cadence still determines how much work must be replayed.
- **Serving recovery follows the state budget**: Preserve authoritative history, then choose reconstruction, checkpointing, or replication for derived state according to the latency contract (table 7.13). A ready replacement provides useful service only when it can also recover the context the request needs.
- **Degradation is a designed mode**: When recovery cannot finish inside the latency budget, a fallback validated in advance (a smaller model, cached results, default features, or shed low-priority load) keeps the service useful, provided monitoring exposes how much quality was traded away.

Fault tolerance makes the cost of coordination a deliberate part of the execution budget. Section B.6.3 works a complete coordination-bound diagnosis, attributing a large run's lost wall time to checkpoints, failures, pipeline bubbles, and scheduling. In the Fleet Law (equation 1.3), checkpoint pauses, detection timeouts, and restarts all land in the synchronization term $T_{\text{sync}}(N)$, and in goodput accounting they are runtime losses. The chapter's mechanisms spend Coordination (C_3) so that a failure does not erase completed Compute, and asynchronous checkpointing and warm restart earn their complexity by moving part of that synchronization cost into T_{overlap} or shrinking it outright. The useful comparison is between time spent preserving and validating state during normal operation and time lost reconstructing trustworthy execution after a failure. A cheaper recovery mechanism is valuable only when it still meets the workload's limits on lost progress, response delay, and degraded quality.

What's Next: From resilience to resource management

Keeping the fleet running despite inevitable hardware failures requires three layers: checkpointing preserves progress, elasticity absorbs node losses, and redundancy protects stateful serving sessions. Each layer also assumed capacity the job cannot grant itself, whether replace-

ment nodes held in reserve, a hot spare kept idle, spot capacity that may vanish, or replicas placed across independent failure domains. Providing that capacity takes allocation decisions, made while many jobs compete for the same fleet, from month-long training runs to hour-long experiments. A fault-tolerant cluster therefore still needs orchestration to decide which jobs run where, when to preempt, and how to share capacity fairly among them.

↪ Admission, placement, preemption, and fair sharing across a multi-tenant fleet are developed in chapter 8.

8

Fleet Orchestration

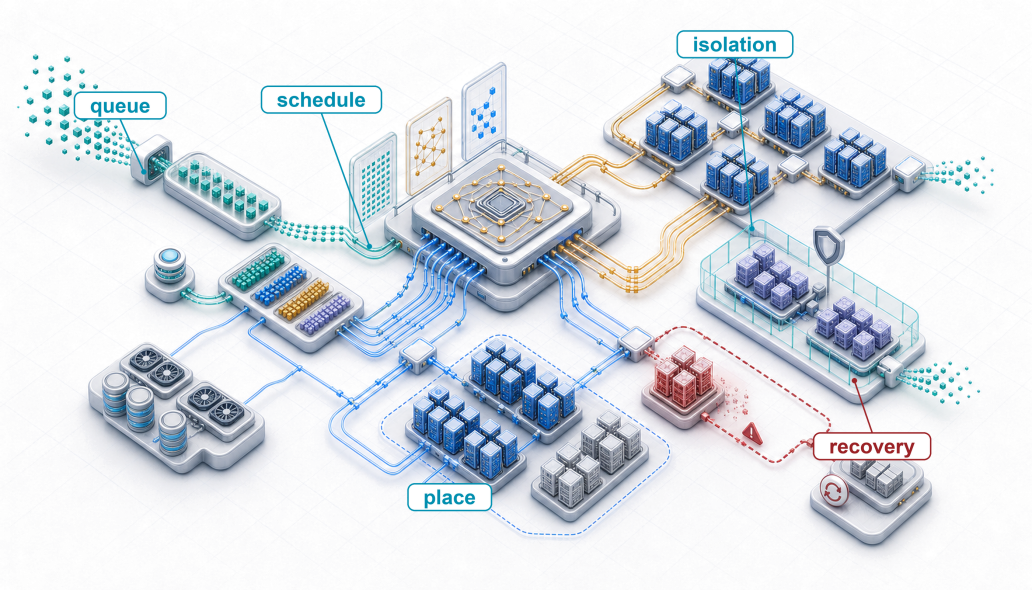

8.1 The Scheduling Problem
8.2 Distributed Scheduling Complexity
8.3 Scheduling Objectives and Their Conflicts
8.4 Core Allocation Constraints
8.5 Orchestration Paradigms
8.6 Topology-Aware Scheduling
8.7 Elastic Scheduling
8.8 Cost Optimization
8.9 Custom ML Schedulers
8.10 Serving Capacity in the Shared Fleet
8.11 Multi-Tenancy and Quotas
8.12 Debugging Cluster Utilization
8.13 Fallacies and Pitfalls
8.14 Summary

Purpose

Why can jobs wait while expensive accelerators sit idle?

A thousand idle accelerators still incur ownership or reservation costs while queued experiments make no progress. More hardware does not resolve that waste when available resources cannot form the groups that pending jobs need. Free devices may be scattered across machines with unsuitable connections, or several jobs may each hold resources while waiting for the rest. On one machine, improving execution usually begins with the work already running. Across a shared fleet, useful progress also depends on which work can start, how long it waits, and what other work its allocation displaces. A short experiment needs a timely result, a long training run needs continuity, and a serving workload needs capacity before requests miss their deadlines. Favoring any one claim without considering the others can leave the fleet busy while delaying the results it exists to produce. Interrupting a job can free capacity, but preserving its state and restoring progress consume time and bandwidth. Waiting for a better placement can accelerate execution while extending completion time. Fair access therefore cannot be judged from device counts alone, and high utilization cannot establish that capacity is being used well. The engineering challenge is to honor these competing claims under changing demand and uncertain resource availability, while keeping the allocation process itself reliable. In C^3 terms, orchestration pays coordination overhead to turn compute and communication capacity into useful progress while preserving fairness and latency headroom.

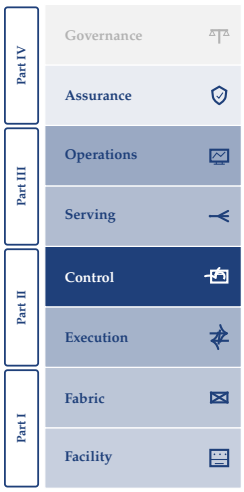

Learning Objectives

- Explain how batch and declarative control planes admit, place, and reconcile allocations under stale state and partial failure
- Apply gang, multi-dimensional fit, and topology constraints before dominant-resource fairness to find allocations a job can run
- Analyze conflicts among throughput, latency, fairness, and cost using slowdown, queueing delay, and ML-specific scheduling signals
- Design topology-aware placements that map tensor-, pipeline-, and data-parallel groups onto NVLink domains and switch tiers
- Design elastic, preemption, and spot policies whose throughput or price gains repay their transition and interruption costs
- Diagnose utilization losses by separating allocated, active, and productive capacity and tracing each gap to its cause
- Evaluate serving-capacity and quota policies that balance serving headroom, team guarantees, and reclaimable borrowed capacity

8.1 The Scheduling Problem

Consider two research teams sharing a 1,000-GPU cluster: Team A submits a 512-GPU job that will run for a month, while Team B submits hundreds of 8-GPU experiments that each run for an hour. If the scheduler continually admits whichever jobs fit, Team B's small experiments can indefinitely block Team A's large training run. The scheduling problem is to resolve this conflict among throughput, fairness, and cluster utilization.

Production traces show that this mix is the normal case rather than a contrived one. In the Kalos cluster that Shanghai AI Laboratory used for language-model development in 2023, 86 percent of 19,907 GPU jobs requested a single GPU, yet together they consumed 0.1 percent of the cluster's GPU time. Jobs of 256 GPUs or more made up 4 percent of submissions and consumed 96 percent of GPU time (Hu et al. 2024), two distributions that figure 8.1 plots side by side. Counted by jobs, the cluster is a stream of small experiments; counted by GPU-hours, it is a few percent of jobs that hold nearly all the accelerators. A scheduler must serve both views at once, which is why fairness, packing, and preemption policies cannot be judged by either count alone.

Mechanisms for synchronizing workers efficiently (chapter 5), partitioning computation across them (chapter 6), and recovering from hardware failures (chapter 7) solve the single-job scaling problem. Each of them assumed that the job already held its accelerators. Fault tolerance assumed more than that. Its recovery paths drew on a spare pool deep enough to replace a drained node and on spot capacity whose loss the job could absorb. Neither exists unless something outside the job grants it, holds it in reserve, or takes it back. That something is orchestration, the policy that decides which jobs run, where they run, when one is preempted for another, and how competing demands share the fleet. An allocation decision must hold across unreliable nodes, balance competing fleet objectives, and satisfy each job's feasibility constraints. Only then can platform control planes, topology-aware placement, elastic and cost-aware policies, workload-aware ML schedulers, and serving, multi-tenant, and diagnostic practices turn raw capacity into useful work. Training and serving draw on the same fleet, so the capacity envelope this chapter allocates is the one both depend on; how a serving system spends its share, request by request, is the subject of chapter 10.

In the fleet stack (shown in the margin), orchestration occupies the Reliability & Control row of the distribution band, but the infrastructure band beneath it constrains its decisions. The scheduler must account for GPU heterogeneity, NVLink topology, rack power limits, and network bisection bandwidth. When scheduling goes wrong, debugging can require examining all four bands to identify whether the root cause lies in infrastructure (hardware constraints), distribution (algorithm limitations), serving (workload mismatch), or governance (policy misconfiguration). Orchestration connects physical infrastructure to distributed algorithms, translating resource requests into placements that respect both hardware topology and organizational policy.

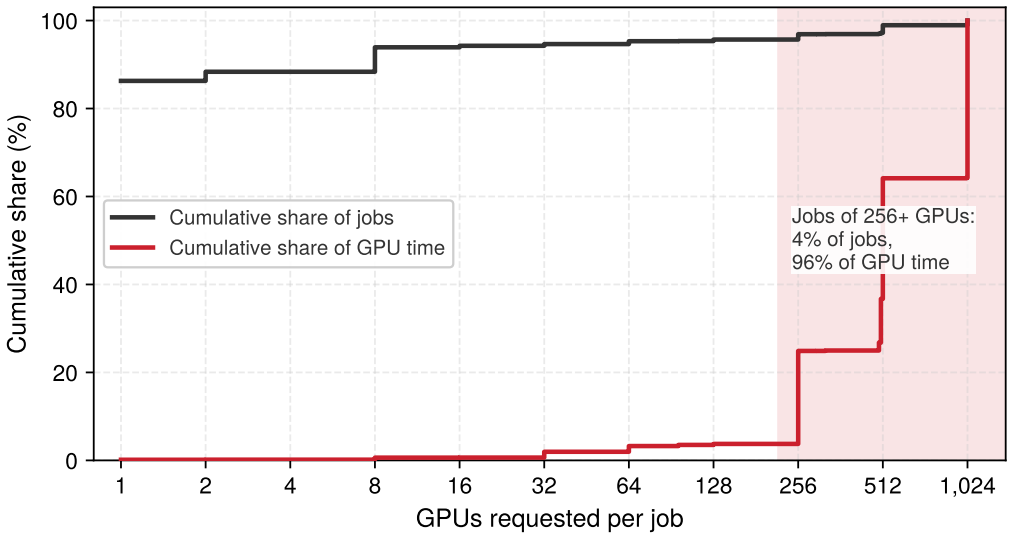

Figure 8.1: Jobs versus GPU Time by Job Size: Cumulative share of GPU jobs (gray) and of GPU time (red) against the number of GPUs each job requested, for the Kalos LLM-development cluster from March to August 2023. Almost every job is small, while almost all GPU time goes to jobs of 256 GPUs or more. Data from the AcmeTrace release of Hu et al. (2024).

Translating those requests into placements grows harder as the fleet and its flagship workloads scale. In a 10,000-GPU cluster, a single pretraining run of a GPT-3-class 175B-parameter model (Brown et al. 2020), the running instance of Archetype A (GPT-4/Llama-3) throughout this book, can hold a 1,024-GPU allocation for weeks. Because the scaling laws of section 1.3 reward larger models trained on more tokens, each new generation of such runs asks for a larger gang held for longer, so the scheduler's hardest requests keep growing. When enough jobs of that class are queued that together they ask for more GPUs than the cluster holds, alongside thousands of smaller experiments and inference workloads competing for the same pool, naive policies break down. A first-come-first-served policy without backfilling (running smaller jobs in the gaps while a large job waits) or preemption can let large jobs occupy the cluster for weeks, delaying smaller experiments that researchers need to iterate quickly. A fair-share policy that divides capacity without respecting each job's minimum resource bundle can leave many partial allocations and no runnable large job.

The economic stakes make these decisions consequential at every scale. A continuously provisioned 10,000-GPU cluster, valued at the chapter's cloud-equivalent rate of $2/GPU-hour, represents $480,000 per day in capacity, whether the GPUs are computing useful work or sitting idle while jobs wait in a queue. If scheduling inefficiencies leave 30 percent of GPUs idle, that translates to $144,000 per day in unused capacity value, or about $52.6M annually. This share counts only the idle capacity that better scheduling can recover, the gap between a poorly scheduled fleet at 50 percent utilization and a well-scheduled one at 80 percent. The poorly scheduled fleet idles more of its GPUs than this, and section 8.8 prices that full idle share.

↩ Archetype A and its 175B-parameter running instance are established in section 1.6.1.1.

Holding useful work per active GPU-hour constant, improving utilization from 50 percent to 80 percent adds 3,000 fully active GPU-equivalents of productive capacity without purchasing additional hardware. That count measures the gain in useful work. Procurement planning asks a different question, namely how much hardware the gain replaces. For that comparison the answer is the gain divided by the 50 percent baseline utilization, or 6,000 raw GPUs, because new GPUs run under the same scheduling policy would also reach only 50 percent utilization. Recovered capacity becomes cash savings only when it avoids purchases or billed resource time. Scheduling can therefore increase the useful work delivered by an existing investment, even when the fleet's immediate expenditure remains unchanged.

Workload semantics, rather than an ML-versus-high-performance computing (HPC) boundary, determine scheduling flexibility. Rigid synchronous jobs in both domains require a fixed set of co-scheduled ranks, whereas malleable jobs can make progress across a declared range of resource

counts. Gang scheduling[1] is the central allocation constraint for a rigid training group: a job configured to require 1,024 GPUs cannot complete its distributed iterations with only half of them. Malleable data-parallel training can instead vary its replica count at explicit reconfiguration boundaries, but tensor- and pipeline-parallel groups usually retain a rigid base allocation. The scheduler must therefore model each job's minimum viable gang, the smallest set of co-scheduled devices on which it can make progress, and its permitted resize points rather than infer rigidity from the workload's field.

If gang scheduling determines which resources must be available together, topology awareness via placement groups[2] determines how efficiently they can communicate. Tensor-parallel groups benefit from placement within a high-bandwidth NVLink domain; a placement across slower links must still meet the job's communication budget. Two allocations of the same size can therefore deliver very different useful throughput.

GPU heterogeneity is the third dimension. Many clusters contain mixtures of A100 and H100 accelerators with different compute throughput and interconnect bandwidth. Their training-throughput ratio depends on the workload, numerical precision, and communication path, so the scheduler needs measured performance and the applicable allocation cost. Memory capacity adds a hard admission test once pools differ in it. The common A100 and H100 variants both carry 80 GB of high-bandwidth memory (HBM), but smaller-memory variants of the same accelerator and older generations often share the fleet. A job whose working set fits only the larger devices cannot run on the rest at any priority. The debugging example in section 8.12 returns to exactly such a mixed fleet. The scheduler must match workloads to appropriate hardware based on actual requirements, not user preferences, while maintaining high utilization across heterogeneous pools. When users request specific GPU types out of habit rather than necessity, the mismatch between requested and required resources can leave entire hardware pools idle while queues overflow on preferred hardware.

Job durations are also uncertain: a training run may converge early and complete in days rather than weeks, hardware failures may abort jobs and return resources at unpredictable times, and hyperparameter searches may terminate unpromising configurations early. In both ML and scientific computing, fixed-step jobs with stable iteration times permit better release-time estimates than jobs whose stopping conditions or execution rates vary. The scheduler must therefore distinguish an estimated completion time from an enforced runtime limit when deciding whether another job can safely use a temporary gap in capacity.

The interaction between these challenges creates combinatorial complexity. A scheduler must simultaneously handle gang scheduling constraints (atomic allocation), heterogeneous resources (capability matching), topology requirements (NVLink locality for tensor parallelism), priority and fairness policies (organizational equity), preemption decisions (which running jobs to interrupt), and cost optimization (spot vs. on-demand placement). Each dimension constrains the others. A topology-optimal placement may violate fairness policies, a cost-optimal spot placement may violate reliability requirements, and a fairness-driven allocation may fragment resources in ways that prevent gang scheduling. The scheduler therefore does more than pack jobs onto GPUs. It must maintain a consistent, economically useful view of a failing, heterogeneous fleet. At that scale, placement stops being only an optimization problem and becomes a distributed systems problem of keeping state coherent across thousands of machines.

8.2 Distributed Scheduling Complexity

A single-machine scheduler works inside one operating-system boundary, where three properties hold closely enough to design around: it observes resource state directly and consistently, its allocation decisions take effect atomically, and failures are nearly binary (the machine is up or it is down). Device faults and delayed cleanup still occur, but they do not hide behind a network. A cluster scheduler surrenders all three properties. It coordinates independent nodes through delayed messages and uncertain observations, so an authoritative allocation decision does not establish that every node has received or completed it. Several distributed systems problems emerge from that gap.

Partial failures pose the first challenge. A node can fail between allocation and job start, creating a gap between the scheduler's decision and its effect. The scheduler may successfully allocate 32

[1] **Gang scheduling**: Formalized by Ousterhout (1982) for parallel systems, where the "Ousterhout matrix" co-schedules related threads across processors in synchronized time slices. In a rigid ML training group, incomplete allocation prevents collective completion, so held resources can remain idle until the missing ranks launch.

[2] **Placement group**: A cloud-native abstraction that requests instances be placed within a high-bandwidth, low-latency network domain. For the ML Fleet, placement groups reduce placement scatter and may expose higher network limits than unconstrained placement. They do not guarantee workload-level bisection bandwidth: topology oversubscription, concurrent traffic, routing, and the collective pattern still determine delivered AllReduce bandwidth.

GPUs across 4 nodes, only to have one node fail before the job launches. The remaining 24 GPUs sit idle while the scheduler detects the failure, updates its state, and re-plans the allocation. Meanwhile, other jobs that could have used those 24 GPUs have already been placed elsewhere, potentially on suboptimal hardware. Before the first training step, recovery must repair allocation and launch state even though no new model progress exists to restore. Once execution begins, resource recovery must also coordinate with the framework's checkpoint and membership mechanisms.

Network partitions create a second problem. The scheduler may lose connectivity to a subset of nodes while those nodes continue operating normally. From the scheduler's perspective, their status is uncertain; from the nodes' perspective, jobs may still be running and producing useful work. Reallocating their GPUs risks double allocation if the partition heals while the old work still runs; waiting for reconnection leaves possibly healthy capacity idle. The partition's duration cannot be known in advance, so any fixed timeout is a guess about network behavior that may prove wrong. A timeout can trigger recovery, but it does not prove that the old execution has stopped. Reusing a GPU therefore requires confirmed cleanup or an enforced revocation of the old allocation. Restarting a job on reachable replacement nodes also requires fencing stale attempts from shared outputs. One common form of fencing advances the allocation's generation number, so that node agents and storage reject any work carrying an older one. The recovery policy can choose how long to wait for reconnection, but each choice must preserve exclusive ownership where the workload requires it.

State inconsistency emerges as a third challenge that compounds the first two. Resource state may differ between the scheduler's view and reality on individual nodes. A GPU the scheduler believes is free may still be cleaning up from a previous job: flushing caches, deallocating device memory, running a zombie process from a failed container, or completing a CUDA driver reset after an error. Conversely, a GPU marked as "in use" may have been freed by a job that terminated between heartbeat intervals. The scheduler therefore operates on a model of cluster state that is always slightly stale. The degree of staleness varies across nodes with heartbeat frequency, network latency, and the complexity of cleanup operations.

↩ Epoch fencing and leases during job recovery are established in section 7.9.2.

Ordering without global time presents the fourth issue. Allocation messages issued before and after a partition can arrive out of order. Logical timestamps, counters that order events without a shared clock, can describe an order, but do not by themselves prevent conflicting ownership. The control plane must serialize conflicting allocation decisions, and the resource manager must reject stale allocation generations before acting on them. Otherwise, an old launch request can revive work after a newer allocation has taken ownership.

These four challenges echo the CAP theorem of section 1.4.2.1, under which a partitioned system cannot guarantee both consistency and availability. CAP is a theorem about distributed data systems, not a literal proof about every scheduler, but the analogy draws attention to which allocation decisions remain safe when nodes cannot communicate. Slurm, an HPC workload manager, and Kubernetes, a container orchestration platform, both need authoritative allocation records and mechanisms that reconcile those records with execution. Optimistic concurrency permits schedulers to plan against potentially stale snapshots, then validate proposed changes before committing them. Rejecting a conflicting proposal costs scheduling work; accepting conflicting ownership would violate a different, essential invariant.

The control-plane tension becomes concrete when the scheduler interacts with the ML framework's own distributed coordination plane during a partition. A scheduler may decide that a subset of nodes is unavailable while workers on those nodes are still blocked inside a collective operation. That creates a brief window of conflicting intent: the scheduler wants to recover the allocation, but the training job has not yet converged on a new membership view. Elastic launch mechanisms such as TorchElastic reduce this ambiguity through rendezvous, the step at which workers agree on group membership and ranks. They use heartbeats and re-rendezvous so the framework can either reform a viable worker group or terminate cleanly, which gives the scheduler a clearer recovery signal than an opaque collective timeout (PyTorch Contributors 2026c, 2026b). This layered coordination, with scheduler-level policy for coarse allocation and framework-level membership for workers, is the characteristic architecture of production distributed training clusters.

8.2.1 Admission, placement, and reconciliation

When a single undifferentiated queue mixes quota checks, hardware binding, and crash recovery, operators cannot tell whether a pending job is blocked by team quota, fragmented topology, or a dead node whose lease, a time-limited ownership grant that lapses unless renewed, has not yet expired. A stale retry can also assign the same accelerators to two jobs at once. A scheduler control plane avoids this ambiguity by separating three decisions. **Admission control** decides whether a job may enter the runnable set under quota, priority, gang, and safety constraints. **Placement** binds an admitted job to a feasible set of nodes while accounting for accelerator capability, topology, power, storage, and network capacity. **Reconciliation** compares desired allocations with observed node and job state, then starts, retries, evicts, or fences work until the two agree. Keeping these stages explicit makes it possible to distinguish admission denial from placement failure, while versioned decisions and resource-level enforcement prevent stale observations from authorizing competing owners.

Control-plane architectures distribute these decisions differently (Hindman et al. 2011; Burns et al. 2016). A centralized scheduler serializes admission and placement against one authoritative state, which simplifies conflict handling but concentrates scheduling latency and failure impact. A two-level design offers resources to framework schedulers, gaining policy specialization at the cost of delayed or rejected offers. A shared-state design lets several schedulers plan concurrently against a common state and resolves conflicts through version checks or optimistic concurrency. In every design, reconciliation closes the loop. Node agents must enforce allocation generations and lease validity so that an old decision cannot revive after a newer placement has taken ownership.

Feasibility also extends beyond free accelerator count. Power-aware admission rejects or delays a placement whose worst-case draw would exceed a rack or facility envelope; relying on average draw risks correlated throttling when colocated jobs enter compute-heavy phases together. I/O-aware placement treats dataset-read and checkpoint-write bandwidth as reservable capacity, so a gang that fits in GPU memory does not start where concurrent input pipelines or checkpoint bursts will saturate the same storage path. These constraints can deliberately leave accelerators idle. The idle capacity is the cost of preserving the power and data-movement budgets required for the admitted jobs to make progress.

8.2.2 Failure as a scheduling input

A scheduler that allocates 10,000 GPUs as if every granted node will stay healthy until the job finishes will see a multi-thousand-GPU training gang stall every few hours waiting for replacement capacity that was already handed out to other workloads. As section 7.2 established, the fleet failure rate is the per-device rate multiplied by the population (equation 7.4), so a cluster of ten thousand GPUs loses a device every few hours (table 7.1). Where fault tolerance turns that failure rate into checkpoints, detection, and elastic recovery for a running job, the scheduler must plan for failure in every allocation decision. The scheduler uses two failure rates, one for each decision. It sizes spare capacity and replacement paths from the GPU-only count of devices it must replace. Its estimate of how often a running job is interrupted, by contrast, uses the shorter full-system interval of section 7.2.2, which adds PCIe and network failure domains, because those faults also stop a synchronous job even when no GPU needs replacing.

The scheduler cannot treat the cluster as a static resource pool where resources, once allocated, remain available until voluntarily released. It must anticipate loss of allocated resources during execution and provide replacement capacity or rescheduling paths. The response depends on the fault and the validity of the surviving state. A transient disruption may permit recovery in place if the execution state remains usable; a persistent hardware failure requires replacement resources. Even a short disruption can force checkpoint recovery if the framework cannot safely continue (Tiwari et al. 2015).

When a node heartbeat is missed, the scheduler faces three recovery policies with different delay and capacity costs. It can wait to see whether the heartbeat resumes, though this delays recovery if the machine has permanently crashed. It can instead begin recovery promptly on replacement resources, provided stale execution is fenced before conflicting work resumes. An intermediate approach marks the node as suspect and prepositions replacement capacity while checking its status; this reduces launch delay if replacement becomes necessary, but holds capacity that could otherwise serve pending work. Whichever policy it picks, reuse of the node's GPUs still waits on the fencing

that network partitions require. Historical failure and recovery measurements set how long to wait and how much replacement capacity to hold.

The recovery-time decomposition of fault tolerance (equation 7.10) shows where scheduler policy intervenes most effectively. Of its detection, restart, load, and warmup terms, the scheduler owns most of T_{restart}, the time to allocate replacement resources and launch processes, and that is the term its spare-capacity policy shortens. Waiting longer on a missed heartbeat lengthens T_{detect} in exchange for avoiding a needless restart, so the wait is worth extending only while a transient fault remains more likely than a permanent one.

These requirements connect the scheduler directly to checkpoint and recovery infrastructure. Scheduling decisions about spare capacity, preemption policies, and elastic training support determine how quickly the system recovers from statistically expected failures. A scheduler that reserves spare capacity or supports elastic recovery can restore throughput much faster than one that requires every failed job to wait for a full fresh allocation and reload. The reserve fraction follows from the measured failure rate and recovery time of the fleet's own workloads. Once the control plane can maintain coherent ownership across admission, placement, reconciliation, and node failures, it still needs an explicit policy objective to decide which competing jobs receive capacity when demand exceeds what the fleet can run at once.

↪ The same decomposition in scheduler-facing phases (detect, reschedule, reload, and replay of work since the last checkpoint) is developed in section E.3.1.

↩ Failure detection, recovery procedures, and spare pools are established in section 7.9.

8.3 Scheduling Objectives and Their Conflicts

Giving a large training gang an unbroken allocation for weeks amortizes launch and data-movement overhead, but leaves short interactive debugging jobs waiting hours in the queue. Preempting the large gang so short jobs start immediately burns cluster capacity on checkpoint saves and restarts. Every scheduling policy trades four competing objectives against one another: throughput, fairness, latency, and cost efficiency. Throughput benefits from allocations that sustain useful work and amortize launch, interruption, and data-movement overhead; a large allocation helps only while the job scales effectively. Maintaining fairness demands an equitable distribution of resources across competing tenants to prevent a single dominant project from monopolizing the fleet. Dominant resource fairness (Ghodsi et al. 2011) formalizes this balance by equalizing each tenant's largest normalized share across GPUs, CPU, and memory. Reducing queue latency helps short, interactive jobs complete promptly, but aggressive preemption can repeatedly interrupt long-running training. Pursuing cost efficiency introduces different costs for different policies: spot capacity risks interruption, off-peak scheduling delays admission, and dense packing can increase interference. Each policy must be evaluated against useful work completed and the delays it imposes.

Fairness requires an outcome metric rather than an input quota rule alone. For job j, **slowdown** $S_j = (W_j + T_j)/T_j$ compares its queueing-plus-running time with its isolated runtime T_j, assuming the running time is unchanged by sharing. A one-hour job that waits nine hours has $S_j = 10$, even if its team eventually receives its nominal GPU share. When contention or resizing changes execution time, the numerator must instead use the observed queueing-plus-running duration while the denominator retains the isolated baseline. A scheduler should report the distribution of slowdown by user and job class, together with the age of the oldest pending feasible job. Dominant-share parity detects unequal resource allocation, slowdown exposes unequal delay, and the maximum feasible wait exposes starvation risk. These metrics separate a policy that is fair in allocated GPU-hours from one that is fair in the time users wait to obtain useful results.

Each of the four objectives, pursued alone, schedules the same pool differently and gives up something the others protect (figure 8.2). Favoring a large unbroken allocation can delay short experiments, while repeatedly admitting short jobs can prevent a large gang from assembling. These conflicts depend on the workload and available capacity; improving one objective does not necessarily worsen every other objective.

Consider a 64-GPU fine-tuning run of the same 175B-parameter model, a far smaller allocation than its 1,024-GPU pretraining run. If it sustains high utilization for days, an uninterrupted allocation amortizes launch overhead across a long execution. Admitting it may nevertheless require draining 64 suitably placed GPUs, delaying smaller experiments. Conversely, repeatedly filling available gaps with short jobs can prevent the large job from assembling its required group. Preemption or resizing can recover flexibility only if the job supports the transition and its cost is acceptable. The policy must therefore compare useful throughput with completion delay for both job classes. In

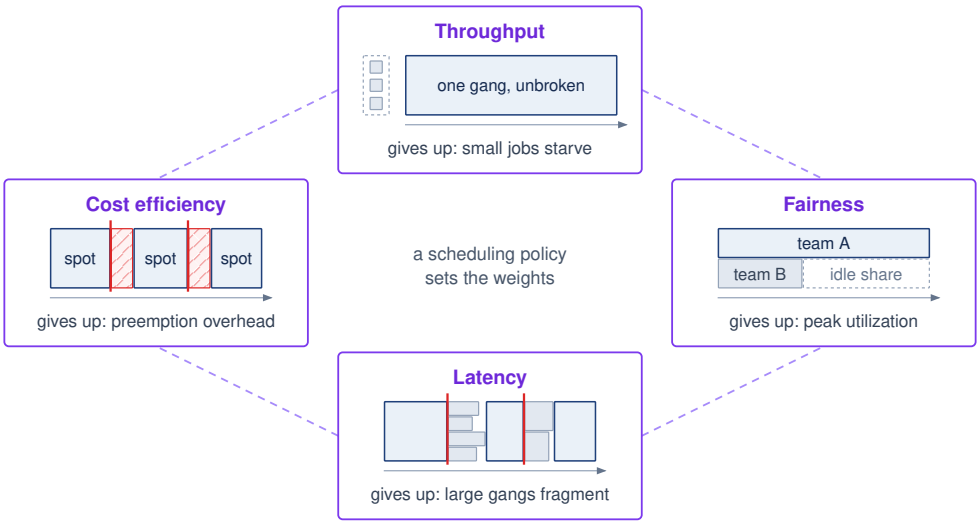

Figure 8.2: Scheduling Trade-Offs across Competing Fleet Objectives: A four-way trade-off diamond (Throughput, Latency, Fairness, Cost Efficiency) illustrates the tensions governing cluster scheduling. Throughput favors long, unbroken runs; Latency favors aggressive preemption for interactive jobs; Resource-share fairness can constrain placement to equalize shares; and Cost Efficiency uses discounted, interruptible capacity. Each vertex sketches a four-GPU pool over time as that objective would schedule it and names a possible cost of favoring it; the conflicts depend on workload and available capacity.

scheduler terminology, those choices often appear as quality-of-service classes: job-priority tiers that decide which objective wins when throughput, latency, fairness, and cost collide.

Queueing adds a further conflict between utilization and latency, because a fleet run close to full load leaves arriving jobs little spare capacity. In a single-server queue with random arrivals, mean waiting time, expressed as a multiple of the mean job duration, is the product of a load factor $\rho_{\text{serv}}/(1-\rho_{\text{serv}})$ and a variability factor $(1+c_s^2)/2$. Here ρ_{serv} is offered load relative to service capacity, and c_s is the coefficient of variation of job duration. The load factor stays small through moderate load and grows without bound as load approaches capacity, so waiting time bends sharply upward near saturation instead of rising in proportion to utilization. The variability factor scales the whole curve. A queue that mixes minute-long debugging jobs with week-long training runs waits several times longer at the same load than a queue of similar-length jobs. Each added point of utilization therefore costs more waiting time in an ML fleet than in a uniform one. Notebook 8.1 states the relation formally and prices both factors.

Napkin Math 8.1: The queuing theory of GPU clusters

To isolate the effect of runtime variability, consider an M/G/1 queue with first-come-first-served service. Kendall notation names the queue by its arrival process, service-time distribution, and server count: M for memoryless (Poisson) arrivals, G for a general service-time distribution, and 1 for a single server. Treating a fleet as one server means modeling its aggregate capacity as a single resource that works through one job at a time, with each job's service time equal to its GPU-hours divided by the fleet's GPU count. That mapping ignores gang packing and parallel admission, so it isolates the variability effect rather than predicting a real fleet's wait. Jobs arrive according to a Poisson process (λ_{arr}), and independent service times follow a general distribution (G) with mean $1/\mu$ and finite standard deviation σ. When offered load is below service capacity, the Pollaczek-Khinchine formula gives the steady-state expected waiting time in the queue (W_q):

$$\frac{W_q}{1/\mu} = \frac{\rho_{\text{serv}}}{1-\rho_{\text{serv}}} \cdot \frac{1+c_s^2}{2}$$

Here, $\rho_{\text{serv}} = \lambda_{\text{arr}}/\mu$ represents offered load relative to the model's service capacity, $1/\mu$ is the average job duration, and $c_s = \sigma\mu$ is the **coefficient of variation** of job duration. The right-hand side is therefore the waiting time expressed as a multiple of the average job duration; a fleet's measured GPU-busy percentage is not automatically this model's load.

An exponential service-time distribution provides a reference case, yielding $c_s \approx 1$. A mixed ML queue can have much greater variability when short debugging jobs measured in minutes share capacity with training runs measured in weeks. The representative $c_s = 3$ calculation in this callout captures that variability without claiming a universal fleet constant. The impact on wait times is multiplicative. For the scheduling argument here, the single-server multiplier is sufficient; section F.1.1 develops an explicit single-server surrogate for batched serving and explains why multi-server, finite-buffer M/G/c/K systems require additional modeling.

At 80 percent utilization ($\rho_{\text{serv}} = 0.8$): with an exponential workload ($c_s = 1$), $W_q = 4\times$ the average job duration. With an illustrative high-variability workload ($c_s = 3$), $W_q = 20\times$ the average job duration. In this M/G/1 model, service-time variability multiplies waiting time at the same offered load. The example does not prescribe a universal GPU-utilization threshold or model gang placement directly. It explains why capacity planning needs the service-time distribution and queue-delay objective as well as average load.

8.4 Core Allocation Constraints

Awarding the 64-GPU fine-tuning run its fair share of the cluster delivers zero progress if those GPUs sit on nodes with exhausted host memory, lack the interconnect its parallel layout requires, or arrive only 32 at a time. A scheduling objective can choose only among allocations that a job can actually run. Fairness must be computed across every resource a tenant consumes, a request must fit every dimension of the node that receives it, a rigid group must arrive whole, and the resources each job holds must not block one another. Each of these constraints narrows the feasible set before any objective selects within it.

8.4.1 Dominant-resource fairness allocation

Equalizing GPU counts alone fails when one tenant's jobs saturate host CPU and memory while another's saturate the network fabric, leaving free GPUs unusable behind exhausted auxiliary resources. Dominant-resource fairness (DRF) makes a multi-resource entitlement comparable by normalizing each tenant's allocation against fleet capacity and taking its largest share (Ghodsi et al. 2011). Consider the deliberately small cluster in table 8.1: 8 GPUs, 32 CPU cores, 256 GB of memory, and 80 Gb/s of network capacity. Tenant A's replica requires 1 GPU, 4 cores, 32 GB, and 10 Gb/s; Tenant B's distributed worker requires 2 GPUs, 2 cores, 16 GB, and 20 Gb/s.

Table 8.1: Worked DRF Allocation: Progressive filling gives Tenant A four one-GPU replicas and Tenant B a pair of two-GPU workers. Both tenants reach a dominant share of 0.50 even though their CPU and memory allocations differ. GPUs and network are jointly saturated; CPU and memory remain available but cannot increase either allocation without exceeding a saturated resource.

Tenant	Allocation	GPU share	CPU share	Memory share	Network share	Dominant share
A	4 replicas	$4/8 = 0.50$	$16/32 = 0.50$	$128/256 = 0.50$	$40/80 = 0.50$	0.50
B	2 workers	$4/8 = 0.50$	$4/32 = 0.125$	$32/256 = 0.125$	$40/80 = 0.50$	0.50

For tenant i, define its dominant share as $d_i = \max_r(a_{i,r}/C_r)$, where $a_{i,r}$ is its allocation of resource r and C_r is fleet capacity for that resource. The calculation equalizes d_i, not every resource share. One A replica raises A's dominant share by $1/8$; one B worker raises B's by $2/8$. Progressive filling, which repeatedly gives one more bundle to the tenant with the smaller dominant share until a resource saturates, therefore grants roughly two A replicas for each B worker until both dominant shares reach 0.50. This is a fair policy target, but it is not yet a feasible placement. If B's two workers form one four-GPU gang, the four GPUs must become available atomically. If tensor parallelism requires one NVLink domain, four free GPUs scattered across nodes do not satisfy the allocation. If the job requires 80 GB HBM accelerators, four incompatible GPUs do not count at all. A production scheduler therefore applies capability, gang, and topology constraints to the feasible set, then uses DRF to choose fairly among allocations that can actually run. The order is not a convenience. Feasibility comes before fairness (principle 15) because a fair share of accelerators on which the job cannot run delivers no progress to the tenant who holds it. When no equal-share placement exists, the scheduler should report the constraint causing the deviation rather than silently presenting an infeasible entitlement as service delivered.

Lighthouse 8.1: Lighthouse archetypes: Placement changes fairness

The same resource share has different meaning across the three lighthouse workloads introduced in section 1.6.1:

- **Archetype A (GPT-4/Llama-3)** needs a rigid, topology-compatible gang. Fractional allocations below its minimum parallel configuration deliver no progress, so fairness must include maximum feasible wait for a contiguous placement rather than GPU-hours alone.
- **Archetype B (DLRM at Scale)** couples accelerator memory to sparse-routing bandwidth. A nominal GPU allocation is infeasible when embedding shards or network capacity cannot follow it, so network and memory must enter the dominant share and placement must protect the serving-latency envelope.
- **Archetype C (Federated MobileNet)** schedules participation rather than data-center GPUs. Eligibility filters for power, connectivity, and device capability can systematically exclude parts of the population, so fairness must be audited through participation and outcome rates across device and user groups, not equal device counts in a round.

The common rule is to define the feasible bundle first and assess fairness over useful progress second. Treating every accelerator or client as interchangeable would conceal the constraint that makes each archetype difficult.

8.4.2 Bin packing

Even when a cluster has dozens of free GPUs in total, a pending 8-GPU job will sit blocked if previous placements scattered those free GPUs in 1- and 2-GPU remnants across partially occupied nodes. Fitting jobs of varying sizes into fixed-capacity nodes without stranding capacity is an instance of **bin packing**.[3] Because general bin packing is NP-hard, no known algorithm finds optimal solutions in polynomial time. Practical heuristics such as first-fit decreasing and best-fit decreasing offer

[3] **Bin packing**: First-fit decreasing has a classical asymptotic $11/9$ approximation ratio for one-dimensional bin packing, with an additive constant in the finite-instance bound. That result does not guarantee the same ratio for ML placement with several resource dimensions, topology constraints, and indivisible gangs. A scheduler must evaluate its heuristic against feasible placements and the stranded capacity it leaves.

tractable choices, but their performance depends on job sizes and placement constraints. How much usable capacity a heuristic recovers in a heterogeneous ML fleet has to be measured against that fleet's job mix, because the classical guarantees cover only one-dimensional packing.

Consider a 64-node cluster with 8 GPUs per node, totaling 512 GPUs. If every job requires 6 GPUs within one node, only one such job fits per node, leaving 2 GPUs unused and reducing effective capacity for this workload to 75 percent. The remaining 2 GPUs per node cannot be combined across nodes into another feasible six-GPU allocation, although smaller compatible jobs could use them. With a mix of 1-GPU, 3-GPU, and 7-GPU jobs, fragmentation depends on the placement and the pending requests; total free capacity can exceed a waiting job's demand while no individual node has enough space. Figure 8.3 maps this stranded capacity across an 8×8 allocation grid. Although 19 slots (30 percent of cluster capacity, dashed white cells) sit idle across nodes 0 through 7, active allocations (J1 through J15) leave every node partially occupied, which blocks a pending 8-GPU job that requires an entire unfragmented node.

ML placement is a multi-dimensional packing problem. Jobs require GPUs together with CPU cores for data preprocessing, host memory for data loading and augmentation pipelines, local SSD storage for dataset caching, and network bandwidth for gradient synchronization. A job requesting 4 GPUs, 32 CPU cores, and 256 GB of RAM may not fit on a node that has 4 free GPUs but only 16 free CPU cores because other jobs have consumed the host CPU for preprocessing. The scheduler must simultaneously satisfy all resource dimensions, and the job fits only if every dimension has sufficient capacity on the selected node. A bottleneck in any one dimension can therefore strand capacity in the others.

Locality constraints further restrict placement beyond simple resource availability. A training job may require tensor-parallel groups to remain within NVLink-connected nodes when its communication budget cannot tolerate the lower per-endpoint bandwidth of the external fabric (450 GB/s per direction intra-node vs. 50 GB/s per port inter-node). On this topology, where NVLink does not connect the nodes, a job requesting "4 GPUs with NVLink connectivity" cannot use 2 GPUs from node A and 2 from node B, even if all 4 GPUs are individually available and the total capacity is sufficient. These topology constraints turn resource counts into a placement problem where which specific resources are available matters as much as how many. A capacity-feasible packing must also pass the job's locality constraints.

Schedulers address fragmentation through several complementary strategies, each targeting a different aspect of the problem. The first, backfill, lets smaller jobs run in the gaps while a large job waits for contiguous resources; section 8.4.3 develops it together with the reservation it must protect and the runtime estimates it depends on. A second strategy, periodic **defragmentation**, migrates or preempts low-priority jobs to consolidate free resources into contiguous blocks, analogous to memory compaction in operating systems. The scheduler identifies nodes where partial allocations leave stranded resources, preempts the jobs occupying those resources, and re-schedules them on nodes where they can be packed more efficiently. The cost includes state transfer and restart, plus replay of work not captured by a valid checkpoint. A graceful checkpoint can preserve recent progress, but still takes time and storage bandwidth. The scheduler must weigh these costs against the benefit of starting large jobs sooner, and low-demand periods reduce the number of competing jobs affected by the move.

A third strategy, **over-subscription**, admits more demand than a capacity dimension can sustain simultaneously, but only where the resource is divisible and the scheduler can bound interference. Average activity of a GPU's streaming multiprocessors (SMs), its compute units, does not create fractional whole GPUs. A job that reserves device memory, requires an exclusive CUDA context, or belongs to a rigid gang still occupies its assigned device while kernels are idle. Statistical multiplexing can be safe for explicitly shareable CPU, network, storage, or partitioned GPU capacity when tenants have measured, weakly correlated peaks and enforceable limits. It is unsafe to convert 70 percent average GPU utilization into a $1.4\times$ device overcommit ratio. Whole-device capacity should instead be recovered through preemption, elastic resizing, or an isolation mechanism such as Multi-Instance GPU (MIG) hardware partitioning or NVIDIA's Multi-Process Service (MPS), which lets several processes' kernels share one device. Either mechanism helps only if the scheduler models its memory, compute, and fault-sharing contract explicitly. Both reappear with the serving workloads they suit best in section 8.10.1.

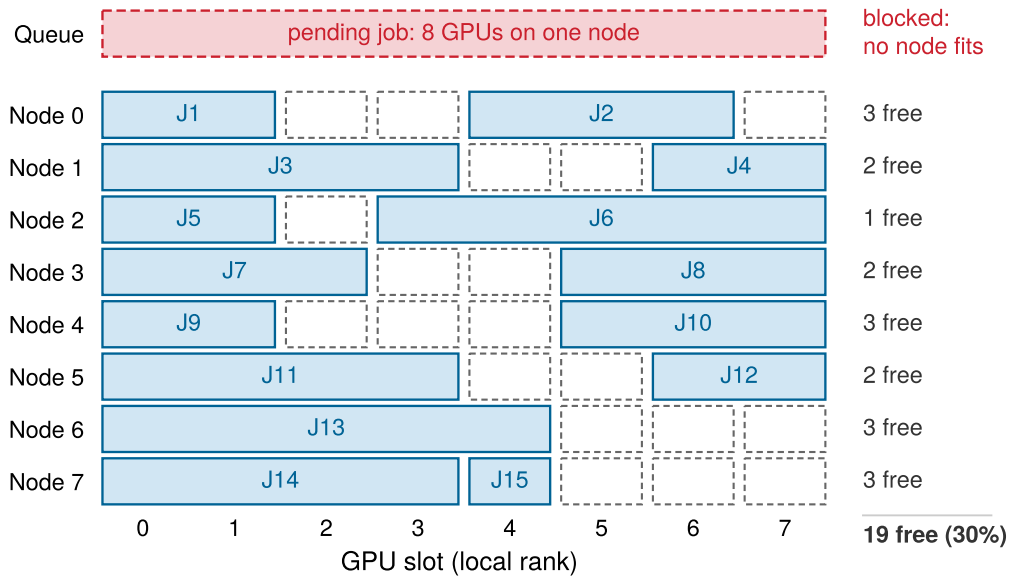

Figure 8.3: The Fragmentation Problem: A snapshot of an 8-node cluster (rows) with 8 GPUs each (columns). Blue blocks mark the slots each active job holds (J1 to J15); dashed white cells are free, and the count at the right of each row gives that node's free slots. Although 19 GPUs, roughly 30 percent of capacity, are free, every node still holds at least one busy slot, so the pending job above the grid, which needs a full node, cannot be placed. Fragmentation, not total capacity, sets what the scheduler can admit.

8.4.3 Gang scheduling

Partial allocation of rigid multi-GPU jobs creates a distributed scheduling failure mode in which two jobs each hold half the resources they need and block each other indefinitely while every allocated GPU sits idle. In the left panel of figure 8.4, jobs J1 and J2 each hold 4 of the 8 GPUs they require, and the whole pool stalls. Atomic gang scheduling (right panel) instead grants J1 all 8 GPUs as one allocation while J2 waits without holding any, which prevents this allocation deadlock.

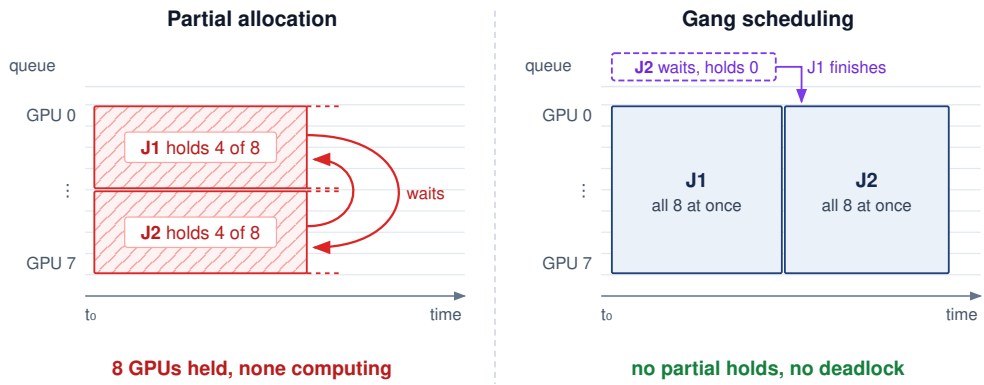

Figure 8.4: The Distributed Deadlock (Hold and Wait): In a non-gang scheduler (left), job J1 acquires half the cluster and waits for the rest, while job J2 acquires the other half and waits for J1 to release its GPUs. Neither can proceed, and 100 percent of the cluster sits idle. Gang scheduling (right) solves this by mandating atomic allocation: a job is either granted its entire N-GPU requirement simultaneously or it waits in the queue without holding any resources.

Definition 8.1: Gang scheduling

Gang scheduling is the ML cluster scheduling policy that atomically allocates the resources required by a job's minimum viable execution group, so a rigid group either receives its complete allocation or waits without retaining a partial one.

1. **Significance**: A rigid synchronous job configured for 1,024 GPUs but holding only 512 GPUs completes zero training steps, not half as many: each collective stalls until every rank in its communication group arrives, so the held GPUs consume power and capital at zero useful output. Without atomic allocation, two such jobs can each acquire half the cluster and deadlock, both at zero progress while consuming the full cluster's operating cost.
2. **Distinction**: Unlike a malleable service or parallel job that can make progress across a resource range, rigid synchronous training has step-function utility: the minimum viable gang or nothing. The distinction is a property of the execution graph, not of whether the workload is labeled ML, cloud, or HPC.
3. **Common pitfall**: A frequent misconception is that priorities or quotas alone solve this. Atomicity must extend to every co-scheduled resource the job needs to make progress (accelerators, high-bandwidth interconnect domains, network capacity for collectives); enforcing it at GPU granularity alone recreates the deadlock one resource layer down.

The partial-allocation pattern is a hold-and-wait deadlock,[4] and its price is the whole allocation for as long as the cycle persists. At the scale of the chapter's reference workload, two 1,024-GPU gangs that have each acquired 900 GPUs from a pool too small to complete either one burn about $1,800 per hour apiece. The limiting case is two jobs that each request the entire cluster, the example OpenAI used to motivate group admission (war story 8.1). On the 10,000-GPU cluster, a deadlock that strands every device wastes $20,000 per hour at the chapter's rate while producing no progress for either job. Measured against the stake computed at the start of the chapter, raising utilization from 50 percent to 80 percent recovers about $144,000 per day in equivalent capacity value. About 7 hours of such a deadlock therefore erase a full day of that gain. Gang scheduling does not create the gain, but it protects it. For a rigid job J requiring N GPUs, the scheduler must grant all N GPUs atomically or leave the job queued without a partial allocation. This removes hold-and-wait for the resources covered by that allocation; it does not prevent later deadlocks in application code or auxiliary resources.

Naive gang scheduling is wasteful in predictable ways. Atomic allocation means the waiting job holds nothing, yet a large gang will never assemble if every freed GPU goes to a smaller job. The scheduler therefore places a **reservation** for the job at the head of the queue, draining capacity toward it by declining to start new work on GPUs that the gang will need. If the cluster has 900 free GPUs and that job requests 1,024, the reservation keeps the 900 GPUs idle until 124 more become available, potentially wasting thousands of GPU-hours. Backfill scheduling addresses this by identifying jobs in the queue that can fit within the 900 available GPUs without delaying the large job's expected start time (figure 8.5), an idea popularized in the ANL/IBM SP scheduling system (Lifka 1995). A 128-GPU job with an estimated runtime of 2 hours is a backfill candidate if the 124 additional GPUs needed by the large job will remain unavailable for at least 2 hours. Protecting the reservation requires an enforced runtime bound or a preemption contract that releases the backfill allocation in time, including any checkpoint and cleanup delay. The estimate identifies an opportunity; the release contract makes the reservation enforceable.

The accuracy of runtime estimates determines backfill effectiveness, and this creates a game-theoretic challenge. If users overestimate runtimes, fewer jobs appear to fit a backfill window. If they underestimate, the scheduler may need to preempt a backfilled job to honor the reservation, paying recovery costs and losing unsaved progress. A policy that kills jobs at their declared limit encourages conservative estimates. Tiresias (Gu et al. 2019), discussed in section 8.9.1, avoids duration estimates when setting priority by using observed resource consumption. That choice addresses ordering under uncertainty, while reservation protection still needs an enforced runtime bound or release contract that says when a backfilled job will give its capacity back.

↩ Ring and tree AllReduce schedules, each of which completes only after every rank's contribution has been reduced, are established in section 5.6.

[4] **Hold-and-wait**: One of the four Coffman conditions (mutual exclusion, hold-and-wait, no preemption, and circular wait), all of which must hold for deadlock to occur (Coffman et al. 1971). Gang scheduling removes hold-and-wait through atomic allocation, which is why it prevents this deadlock by construction; the price is reduced packing flexibility.

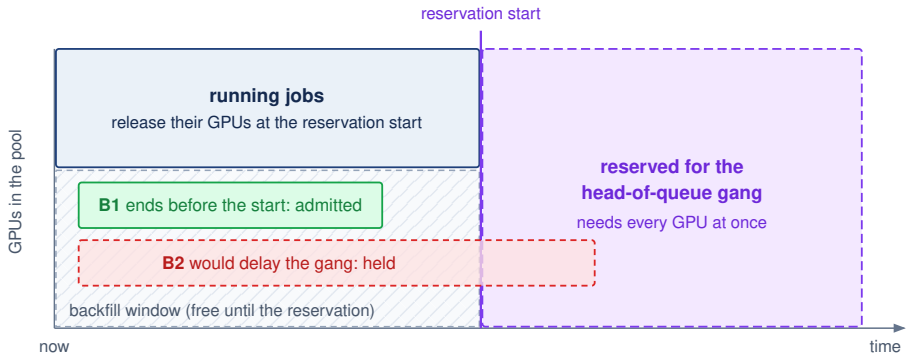

Figure 8.5: Backfill under a Reservation: The head-of-queue gang holds a reservation that begins when the running jobs release their GPUs. Until then, the free GPUs form a backfill window. A short job that ends before the reservation start runs in that window, while a job whose runtime bound would carry it past the start is held, because admitting it would delay the gang. The schematic carries no measured values.

8.4.4 Deadlock prevention and detection

Gang scheduling eliminates the hold-and-wait condition for the atomically allocated bundle, but it does not immunize the cluster against all forms of resource contention. Later requests for auxiliary resources can still create a cycle of dependencies. Priority rules can also prevent a resource holder from making progress without creating a cycle. Distinguishing these cases matters because breaking a deadlock and providing a starved job with enough resources to exit require different interventions.

Priority inversion occurs when a high-priority job waits on a low-priority resource holder whose progress other work delays. Consider the 64-GPU fine-tuning run of the 175B-parameter model, submitted at high priority. The scheduler has reserved 60 available GPUs, but the remaining 4 are held by a low-priority data processing job. Normally, the scheduler would preempt the low-priority job to satisfy the high-priority request. However, if a stream of medium-priority development jobs saturates the cluster's CPU or network bandwidth, starving the low-priority job of the resources it needs to checkpoint and exit, the low-priority job stalls. Its graceful release path cannot complete. The high-priority job waits for the low-priority job, which is effectively blocked by the medium-priority jobs. The 60-GPU idle block persists until the exit path receives resources or an enforced termination policy reclaims the remaining GPUs, potentially losing uncheckpointed work.

Priority inversion: the high-priority job waits on a low-priority holder that medium-priority work starves.

Definition 8.2: Priority inversion

Priority inversion is an ML cluster scheduling pathology in which a high-priority task is forced to wait for a lower-priority task to release a shared resource.

1. **Significance**: It ties the progress of a high-priority job, and the accelerators reserved for it, to the progress of the lowest-priority job in the wait-for chain. In ML clusters, this typically occurs when a low-priority job holding GPUs is starved of auxiliary resources (for example, bandwidth for checkpointing), preventing it from finishing and releasing the accelerators needed by high-priority workloads.
2. **Distinction**: Unlike standard queuing (where tasks wait their turn), priority inversion involves an active blockage: the high-priority task is ready to run but is transitively dependent on a task that the scheduler does not prioritize.
3. **Common pitfall**: A frequent misconception is that strict priority levels solve this. Raising the waiting job's priority does not supply the CPU or bandwidth that its lower-priority resource holder needs to exit; the preemption policy must protect that release path.

Prevention and detection address different parts of this problem. Atomic allocation prevents hold-and-wait for the resources included in the gang, while reserving checkpoint bandwidth allows a preempted job to complete its exit path. Timeouts instead bound how long the system waits before investigating or reclaiming an allocation; they do not prevent deadlock states from arising. Monitoring can trigger recovery when progress stops, but safe reclamation still requires enforcing the new ownership decision.

Detection, unlike exact placement search, is computationally cheap. Before feasibility filtering, a cluster with 10,000 GPUs and 500 pending jobs can have more than 10^{15} candidate allocations, which puts exhaustive allocation search out of reach. Detection instead works on a **wait-for graph**, a directed graph with one vertex per job and an edge from job A to job B whenever A waits for a resource that B holds. A cycle identifies a potential deadlock. It establishes deadlock when every edge is an unavoidable blocking dependency and no alternative resource allocation can let a member proceed. Finding a cycle in an already constructed graph takes time linear in its vertices and edges. The difficulty lies in construction. Resource state changes continuously across a distributed control plane, so the scheduler rarely holds a complete, current graph, and a cycle detected on stale state may already have resolved while a real one goes unseen.

Large-scale schedulers therefore supplement wait-for detection with mechanisms that preserve liveness under incomplete state. A **lease-based resource** gives every allocation a time-to-live; expiry authorizes reclamation only when the resource manager enforces revocation and cleanup before reuse. **Progress monitoring** checks advancement through signals such as completed-step counters, using GPU utilization and heartbeats as supporting observations rather than proof of useful work. When the scheduler must preempt a lower-priority job, **holistic preemption** reserves the CPU, network, and storage bandwidth needed for that job to checkpoint and exit cleanly; otherwise, the victim can be too starved to release the accelerators. False failure suspicions can cost useful work, but must not bypass the ownership checks required for safe recovery.

8.4.5 Heterogeneous gang scheduling

In preference-training pipelines such as reinforcement learning from human feedback (RLHF), allocating the policy-training GPUs while the rollout-generation or reward-scoring workers sit queued stalls the workflow just as completely as a missing AllReduce rank. Homogeneous gang scheduling allocates N identical workers executing the same compute graph. Heterogeneous workflows instead combine trainers, evaluators, and inference-style workers that need different GPU counts, memory footprints, and lifetimes, yet still must run together as a single job.

↩ The models that alignment training co-locates, and their unequal resource demands, are established in section 6.7.

> **Systems Perspective 8.1: Routing over synchrony**
>
> Heterogeneous gang scheduling extends atomic allocation to the roles that must coexist for a workflow to progress. A viable group might pair high-HBM nodes for memory-bound generation with compute-oriented nodes for policy updates. The orchestration problem then shifts from packing a homogeneous ring toward routing data. Rather than reducing gradients among identical workers, the job streams rollouts from generation workers to the workers that compute policy updates, so placement must also support the data exchanged between those roles. Rollout and training stages may exchange work in batches or through a buffered pipeline; their measured production and consumption rates determine where queues grow or accelerators wait. The scheduler needs a minimum viable resource bundle and a supported scaling policy for each role, rather than an assumption that all workers are interchangeable or perfectly balanced.

Homogeneous and heterogeneous jobs share an allocation requirement: the scheduler must assemble a viable group before expecting progress. Backfill can use capacity while that group waits, provided runtime estimates or preemption contracts protect the reservation. Enforcing these policies requires a control plane that represents ownership and reconciles allocations after failures.

8.5 Orchestration Paradigms

A pending gang and a queue of small backfill jobs need both a policy for sharing capacity and a control plane that can enforce it. The policies just developed apply across platforms; the orchestration paradigm determines how the platform represents allocations and recovers when observed state diverges from them. Slurm, shaped by HPC batch systems, favors explicit resource requests and predictable allocation. Kubernetes, shaped by cloud-native service management, favors declared desired state and continuous reconciliation. The choice is therefore not Slurm vs. Kubernetes as products; it is which control model best matches the workload mix, failure behavior, and operational culture of the fleet.

The two platforms embody two control models (figure 8.6). Slurm follows an imperative batch model, in which the user states the resources a job needs and the central slurmctld dispatches that allocation through per-node slurmd daemons. Kubernetes follows a declarative model, in which the user states what should be running and controllers reconcile observed state toward that declaration through the API server and node agents. Both platforms accept explicit resource requirements and track execution state, so the useful distinction is which lifecycle the platform manages directly: an allocated batch job or a collection of workloads maintained by controllers. That distinction shapes how ML workloads are admitted, monitored, and recovered from failure, and distributed training still needs framework-level coordination under either model.

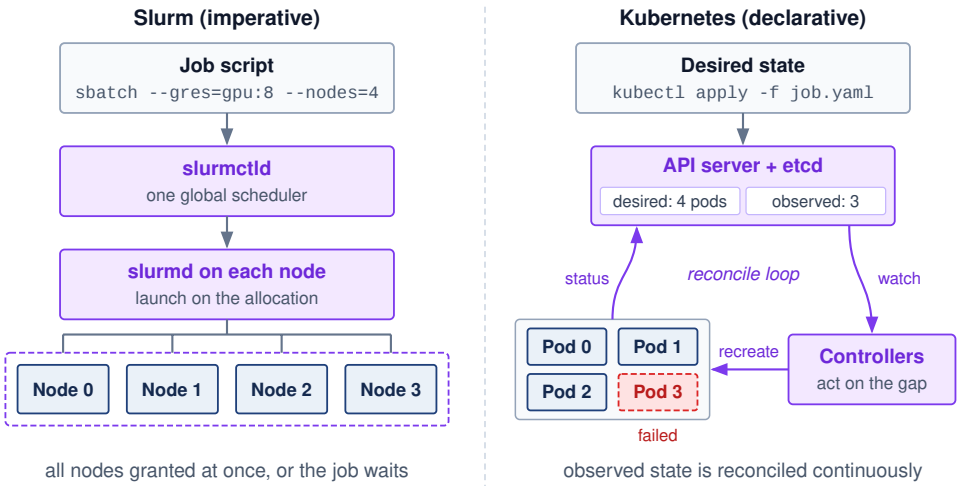

Figure 8.6: Architectural Paradigms: Slurm vs. Kubernetes: Slurm (left) uses an imperative batch model where users submit explicit resource requests to a centralized controller (`slurmctld`) providing deterministic atomic allocation. Kubernetes (right) uses a declarative service model where controllers continuously reconcile observed cluster state in `etcd` against desired state specifications.

8.5.1 Slurm and HPC scheduling heritage

When a multi-node job needs an exclusive, non-overlapping block of accelerators for days at a time without partial-allocation deadlocks or conflicting node ownership, a centralized batch allocator is the most direct control plane. The Slurm Workload Manager[5] originated at Lawrence Livermore National Laboratory to solve that problem for Linux supercomputers (Yoo et al. 2003): long-running jobs, expensive shared hardware, and users who can specify resource needs in advance. In ML infrastructure, the same batch model maps to GPU training through partitions for heterogeneous accelerator pools and predictable atomic allocations for long-running distributed jobs.

Slurm's batch submission model makes resource requirements explicit. Users submit job scripts, and the scheduler allocates resources within eligible partitions according to policy and availability. A submission of `sbatch --gres=gpu:8 --nodes=4` requests 8 GPUs on each of 4 nodes, a 32-GPU allocation. The request supplies a capacity contract, but useful execution still depends on the selected

[5] **Slurm Workload Manager**: Originally "SLURM," an acronym for Simple Linux Utility for Resource Management, the project dropped the acronym capitalization in 2012. Developed at Lawrence Livermore National Laboratory beginning in 2002, Slurm's centralized controller (`slurmctld`) maintains a single authoritative view of resource allocations. At very large scale, operators may partition fleets or federate schedulers to keep scheduling latency, policy isolation, and failure domains manageable.

hardware and the application's launch configuration. Overstating requirements strands resources; understating them can cause memory failures or poor performance.

Slurm's architecture consists of a central controller daemon (`slurmctld`) that maintains the global view of cluster state and makes scheduling decisions, and per-node daemons (`slurmd`) that manage local resources and execute jobs. This centralized architecture provides the authoritative allocation model discussed in section 8.2: the controller is the source of truth for resource assignments, reducing the risk of conflicting allocations. The cost is a potential single point of failure (addressed through active-passive failover) and a scheduling-throughput limit. As fleets grow, organizations may partition the cluster or federate schedulers to bound scheduling latency and failure impact; within a cluster, a typical ML configuration defines Slurm partitions by accelerator type and interconnect (table 8.2).

Table 8.2: Slurm Partition Configuration: Partitions organize heterogeneous accelerators into logical pools matched to workload characteristics. NVLink-connected partitions support tensor parallelism, while PCIe partitions serve workloads that rely primarily on data parallelism. Separating inference and debug partitions prevents experimental workloads from impacting production serving.

Partition	GPUs/node	Interconnect	Typical use
dgx-a100	8× A100	NVLink + IB NDR	Large LLM training
a100-pcie	4× A100	PCIe + IB HDR	Medium training
inference	2× A10G	Ethernet	Model serving
debug	1× V100	Ethernet	Development

Within a partition, a request specifies both the number of GPUs a job needs and the **allocation shape**: how those GPUs are distributed across nodes. In Slurm, `--gres=gpu:N` or `--gpus-per-node=N` requests a per-node count, while `--gpus=N` requests a job-wide total. Whether the scheduler may split a node between jobs is a separate consumable-resource policy (Slurm's `select/cons_tres` plugin), and a node-exclusive policy strands the unused devices in the six-on-eight example of section 8.4.2. Shape decides communication. In an eight-GPU-per-node pool, a 16-GPU request that also fixes eight GPUs per node (`--gpus-per-node=8`) yields two fully populated nodes, whose NVLink domains follow from the node hardware rather than the count. A request for 16 GPUs with no per-node constraint gives the scheduler more freedom, so it can start sooner on partially occupied nodes, but it may spread the group across several of them. Whether that flexibility is useful depends on the parallelism configuration and its communication budget. Tensor-parallel groups with frequent exchanges may justify compact placement, while data-parallel groups with well-overlapped communication may tolerate a wider allocation.

The **fair-share scheduling** mechanism uses past consumption to influence future priority under contention. An illustrative heuristic combines base priority, the user's target fair share (F_{target}), recent actual consumption (F_{actual}), and a small stabilizing constant (ϵ_{stab}); it illustrates the feedback without reproducing Slurm's multifactor priority calculation:

$$\text{Priority}_{\text{effective}} = \text{Priority}_{\text{base}} \times \frac{F_{\text{target}}}{F_{\text{actual}} + \epsilon_{\text{stab}}} \tag{8.1}$$

With base priority and target share fixed, equation 8.1 lowers priority as recent consumption increases. Doubling consumption approximately halves this illustrative priority when the stabilizing constant is negligible. Actual admission also depends on other priority factors and feasible resources; low recent usage does not guarantee an immediate start.

The time decay of usage history determines how quickly the system "forgives" past heavy usage. A half-life of one week means that a researcher's heavy usage from two weeks ago contributes only 25 percent to their current usage calculation. Short half-lives make the policy more responsive to recent changes but can encourage alternating bursts of usage. Long half-lives retain the effect of a completed large project when its researchers return with smaller jobs. The useful decay window depends on the fleet's job durations and the period over which the organization intends to balance access. Usage itself need not count every device alike. **Trackable resources (TRES)** let accounting

charge different accelerator classes different billing weights in fair-share usage, an administrative valuation that tracks economic cost only as far as the chosen weights do.

Preemption policies enable high-priority jobs to reclaim resources from running workloads, and for ML training the reclaim must leave time to save state. Slurm expresses that time as a grace period. It signals the selected job (`SIGTERM`), waits the configured `GraceTime`, and only then kills it (`SIGKILL`); under `PreemptMode=REQUEUE` it can also requeue an eligible job to run again later. Neither the signal nor the requeue saves model state. The application must handle the signal, publish a valid checkpoint within the grace period, and restore it on restart. If termination precedes completion of that save, recovery uses the last valid checkpoint and replays subsequent work. Preemption therefore buys access for high-priority work with the victim's replay and restart time, a price that section 8.8.1 develops.

↩ Checkpointing for progress preservation, including the checkpoint interval and coordination patterns, is established in section 7.8.

Several Slurm features expose ML structure to the batch system. A hyperparameter sweep is the opposite of a gang. Job arrays submit many trials as one record (for example, `--array=0-99`), and each trial reads its index from the `SLURM_ARRAY_TASK_ID` environment variable to select its configuration. Slurm schedules the elements independently, so individual trials become ideal backfill candidates that fill small gaps in capacity, while the array keeps submission and management overhead low without making the trials one indivisible gang.

Heterogeneous jobs carry the heterogeneous gang of section 8.4.5 into a batch allocation. One job can contain components with distinct resource requests, admitted together. The 64-GPU fine-tuning run of the 175B model, for example, could request 8 GPU nodes for training and 2 CPU-only nodes for on-the-fly tokenization as a single allocation. It could then launch training on the GPU component and data workers on the CPU component (`srun --het-group=0` and `--het-group=1`). This separation helps only if the data workers and transfer path sustain the training stage's consumption rate; a heterogeneous allocation exposes that choice but does not itself balance the pipeline.

At the boundaries of each allocation, prolog and epilog scripts turn job transitions into health checks. A prolog tests reported topology, ECC counters, and link status on each allocated node and drains a failing node before launch. An epilog removes the job's orphan processes and shared-memory segments. Draining a failing node this way applies the quarantine practice of section 8.6.4 at job boundaries, keeping sensitive work off a node whose health evidence shows degradation. Passing the prolog reduces exposure to known faults but does not establish the bandwidth the job will receive.

Slurm's strengths for ML training include explicit batch allocations, configurable fair-share policies, and integration with MPI-based distributed training frameworks. A mixed training-serving fleet also needs a service lifecycle that maintains ready replicas as demand and failures change. Inside a batch scheduler, adding or removing an inference replica means submitting or canceling a job, a far heavier operation than a controller adjusting a replica count, and a failed replica stays down until something resubmits it. Operators must integrate replica health, replacement, and demand-driven scaling with the batch scheduler, or assign those responsibilities to a service orchestration platform.

8.5.2 Kubernetes declarative container orchestration

When a cluster runs always-on inference services alongside training jobs, managing serving replicas as batch scripts forces operators to rebuild health checks, rolling updates, and replica replacement by hand. Kubernetes addresses this mixed workload by managing both services and batch jobs through a **declarative model** that expresses a workload's desired state and uses control loops to reconcile observations with that state (Burns et al. 2016; Cloud Native Computing Foundation 2024). The unit it places is the **pod**, one or more containers that share a node and a network identity, and its default scheduler places pods one at a time. The name collides with the hardware pod of section 2.7, the many racks joined under one fabric. A Kubernetes pod is a software unit that never spans more than one node, so a 64-pod training job is 64 container groups placed onto nodes, all of which can sit inside one hardware pod. Controllers above the scheduler give groups of pods a purpose: a Deployment keeps a set number of identical replica pods running, while a Job runs pods until they complete, retrying failures under a stated policy. The workload controller therefore determines what completion, replacement, or continued availability means.

Failure handling reveals the importance of the workload contract. Slurm can detect node failures and requeue eligible jobs when configured to do so. Kubernetes controllers maintain workload-specific goals, such as a Deployment's replica count or a Job's completion and retry policy. A replacement pod still has to be placed, become ready, and rejoin the distributed training run before progress resumes. Training operators can coordinate worker creation, while the framework remains responsible for valid state and group recovery. Both platforms therefore require an explicit agreement between allocation, process lifecycle, and application recovery.

Kubernetes' control loop architecture uses the controller pattern:[6] controllers observe and update API objects through the API server, which persists cluster state in etcd, a replicated key-value store. Those objects contain desired state and reported observations, not instantaneous physical truth. The pattern repeats at every level: the Deployment controller reconciles the number of pod replicas, the scheduler binds pending pods to nodes, the kubelet (the agent on each node) keeps its assigned containers running, and the node controller marks nodes that stop reporting. Each component acts independently and communicates only through API objects, which makes the system resilient to the failure of any single controller. The price of that independence is that their interactions require clear ownership and retry behavior, particularly when a failed node's last reported state no longer describes its execution.

Kubernetes needs an explicit resource representation before it can place GPU workloads. In the device-plugin model, **device plugins** expose accelerators as schedulable extended resources. The NVIDIA device plugin registers GPUs with the kubelet, enabling pod specifications that request GPU resources declaratively through the standard Kubernetes resource model. A pod then asks for whole GPUs by naming the `nvidia.com/gpu` resource in its limits (listing 8.1).

Listing 8.1: Kubernetes GPU Allocation: Pod resource specification requesting GPUs through the NVIDIA device plugin. The `nvidia.com/gpu` resource name follows Kubernetes extended resource conventions, where the domain prefix identifies the device plugin vendor. This declarative syntax enables portable GPU workload definitions across any Kubernetes cluster with the NVIDIA device plugin installed.

```
# Kubernetes pod resource specification for GPU allocation
resources:
  limits:
    nvidia.com/gpu: 4  # Request exactly 4 GPUs for this pod
  requests:
    cpu: "32"
    memory: "256Gi"
```

When the device plugin exposes exclusive whole GPUs, a small model serving occasional requests can reserve a device while using only a fraction of its compute and memory capacity. Whole-device allocation suits workloads that need the capacity or isolation, but can strand resources for smaller workloads. The resource unit exposed by the plugin therefore matters, because sharing or partitioning changes what the scheduler can allocate and what isolation each allocation receives. Device plugins can also expose hardware partitions, such as MIG instances, as schedulable units, a choice that section 8.10.1 develops with the serving workloads it suits.

Kubernetes' default scheduler places pods independently, so it does not supply the group admission that section 8.4.3 requires. Under that default, a 64-pod training job receives 64 separate placement decisions, with no guarantee that all will succeed. If the cluster has resources for 60 pods but not 64, those 60 can occupy capacity while the remaining 4 wait. The platform must therefore configure a group-level admission and placement contract for a rigid job.

The mechanism that closes this gap is group admission. A **PodGroup** declares a group's minimum runnable membership. The scheduler tentatively reserves a placement for each member as it finds one but binds none of them to nodes until the minimum membership is placed, and it releases every reservation if the attempt fails or times out. Admission and placement thereby become two decisions, with the group admitted only when placement has succeeded for enough of it. The Volcano[7] batch scheduler and the Coscheduling scheduler plugin both implement this pattern. Group admission does not make physical launches simultaneous; nodes can still fail or start at different times, so the framework must establish the communication group before training proceeds.

6 **Controller pattern (reconciliation loop)**: A "level-triggered" design reconciles the current desired and observed states rather than depending on every intermediate event. A restarted controller can resume from durable state if its actions tolerate retries and its observations expose unfinished work. Multiple controllers still need distinct responsibilities so that one controller's correction does not repeatedly undo another's.

7 **Volcano**: Volcano provides a scheduler with batch and gang-scheduling mechanisms. Pods select it through their scheduler name, so it can coexist with Kubernetes' default scheduler. The operational contract must account for interactions between schedulers sharing the same nodes; installing Volcano does not automatically transfer every workload's scheduling to it.

War Story 8.1: OpenAI's gang-scheduling example (2021)

Context: OpenAI scaled Kubernetes clusters to 7,500 nodes for workloads including GPT-3, CLIP, and DALL-E. Its account explains that MPI-coordinated optimizer groups needed all members scheduled before training could progress (Sigler 2021).

Mechanism: The account illustrates the risk with two experiments each requesting the entire cluster. Incremental pod placement could allocate half the capacity to each, leaving neither with a complete group.

Impact: In that example, both jobs hold capacity without making progress. It demonstrates the deadlock mechanism rather than documenting a particular pair of stalled production runs.

Fix: OpenAI reports adopting the coscheduling plugin to add group-aware admission to the Kubernetes scheduler.

Systems lesson: Distributed ML scheduling must encode the job's execution semantics. For MPI-style training, "some pods scheduled" is not partial progress; it is wasted allocation. Production schedulers need gang scheduling, quota controls, and workload-aware policies rather than generic pod placement alone.

Priority classes control scheduling precedence and preemption eligibility. A fleet might give inference higher priority than training to help protect serving service-level objectives (SLOs), but priority cannot replace ready capacity during a demand spike. With `preemptionPolicy: PreemptLowerPriority`, a pending pod can displace lower-priority pods when their removal makes placement feasible. With `preemptionPolicy: Never`, it retains queue precedence without initiating preemption and remains subject to higher-priority preemption. Termination, cleanup, and replacement startup take time. Training applications also need a checkpoint-aware shutdown path that fits the allowed grace period; assigning priorities does not create that path. The policy must balance serving readiness against repeated interruption of training.

Kueue[8] makes admission a separate control decision. It manages when workloads may consume quota, while a compatible scheduler binds their pods to nodes. The admission policy and placement policy must agree on resource types, group requirements, and topology. Separating these responsibilities lets each layer run its own policy, at the price of keeping the two consistent. Kueue represents each team or project as a queue with a configured quota (a ClusterQueue) and groups queues into cohorts whose members may lend one another unused quota, under the borrowing and reclaim contract of section 8.11.2.

[8] **Kueue**: Kueue separates workload admission from pod placement. Its all-or-nothing mechanisms include waiting for admitted pods to become ready and requeuing workloads that exceed the configured timeout. Aggregate quota admission alone does not prove that a gang fits the physical topology; the guarantee depends on the configured readiness, topology, and placement mechanisms.

8.5.3 Kubernetes ecosystem for ML training

Group admission handles one constraint that microservice orchestration can leave implicit but distributed training cannot. Four other such constraints remain. Workers must launch knowing their peers, collective traffic must avoid network layers built for request traffic, checkpoint writes need bandwidth rather than only storage capacity, and progress must be observable below the container. Operators and plugins add each constraint to the platform; the mechanism matters more than the product that happens to implement it.

Peer discovery is the job of the **training operator ecosystem**, which extends Kubernetes with job-specific resource types (custom resource definitions). In the Kubeflow Training Operator, a `PyTorchJob` declares the number of workers, the container image, and resource requirements. The operator creates the pods and injects the rendezvous address (`MASTER_ADDR` and `MASTER_PORT`), the group size, and each worker's rank (`WORLD_SIZE` and `RANK`), which the framework uses to form its communication group. If a worker pod fails, the operator can detect the deviation and create a replacement pod to restore the desired replica count, but training-state recovery still depends on the framework's checkpoint and rendezvous logic.

Collective traffic is the second constraint. The pod overlay network that the Container Network Interface (CNI) configures, designed for microservice requests, routes packets through software switching on the host CPU, which adds latency and caps bandwidth when it sits on the critical gradient-synchronization path. For bandwidth-sensitive training, **host networking** (`hostNetwork: true`) places pods in the host's network namespace, bypassing the ordinary pod overlay path. It

does not itself allocate remote direct memory access (RDMA) devices or grant permission to use them; device exposure, resource requests, drivers, and access controls require separate configuration. Operators such as NVIDIA's Network Operator automate that exposure through RDMA device plugins and GPUDirect RDMA configuration, letting containerized workloads approach bare-metal interconnect performance once the rest of the stack is configured.

Checkpoint bandwidth is the third constraint, because storage integration must account for the checkpoint storm of definition 4.5, in which many workers write large amounts of state together. Kubernetes PersistentVolumeClaims (PVCs) request capacity on a configured storage system but do not reserve checkpoint throughput. A 1,024-GPU training job writing the 2.45 TB resumable checkpoint of the 175B model needs a measured write budget shared with other workloads. Backend bandwidth controls, client throttling, or staggered writes can limit interference; the platform must verify which controls the storage implementation enforces.

Visibility into progress is the fourth. The **observability stack** for training clusters combines general-purpose cluster monitoring (commonly Prometheus and Grafana) with accelerator telemetry such as the Data Center GPU Manager (DCGM) exporter, which reports utilization, memory bandwidth, temperature, and ECC errors. These observations locate symptoms. Correlating them with step progress, input waits, and communication traces is what distinguishes useful computation from stalls and identifies which dependency limits a training job.

The 1,024-GPU pretraining run for our 175B model illustrates how these responsibilities fit together. One possible deployment defines a `PyTorchJob` with 128 worker pods, each requesting 8 GPUs. Host networking removes the pod overlay from the gradient path, and the Network Operator exposes each node's InfiniBand adapters to the workers through RDMA device plugins. A PVC backed by Lustre, a parallel file system, provides access to the 6 TB tokenized training corpus and a checkpoint path whose sustained bandwidth must be validated. DCGM metrics stream to Prometheus, supporting alerts on uncorrectable ECC errors and deviations from workload-specific communication baselines.

↪ Fleet telemetry across device, network, and facility levels, and the correlation that narrows a failure domain, is developed in section 12.6.1.

8.5.4 Paradigm trade-offs and selection criteria

Selecting a control plane requires matching its allocation and recovery contract to the workload mix. Rigid training groups need reliable atomic placement and checkpoint-aware interruption; serving replicas need readiness checks and continuous recovery. Existing operational expertise and the cost of running a second control plane also shape the choice. A fleet dedicated to long, homogeneous training runs that use bare-metal collectives, operated by a team fluent in batch scheduling, gains most from Slurm's authoritative allocation and mature fair-share. A fleet that mixes training with always-on serving, and already runs cloud-native deployment and observability tooling, gains more from Kubernetes with a gang-admission layer added for its rigid jobs. The Slurm vs. Kubernetes trade-off organizes around the dimensions most relevant to ML workloads (table 8.3).

Table 8.3: Slurm vs. Kubernetes for ML Workloads: The two paradigms reflect different design philosophies optimized for different workload characteristics. Most production ML platforms use elements of both, with Slurm managing dedicated training clusters and Kubernetes managing inference serving and pipeline orchestration.

Dimension	Slurm	Kubernetes
Scheduling model	Imperative: user specifies exact resources and scheduler allocates them	Declarative: user specifies desired state, system reconciles continuously
Failure handling	Configuration-dependent requeue and checkpoint workflows	Self-healing through control loops
Gang/atomic multi-node start	Atomic job allocation plus backfill; Slurm "gang scheduling" specifically refers to time-slicing/preemption	Production clusters commonly use PodGroup-based support such as Volcano, Coscheduling, Kueue, or upstream alpha gang scheduling
Fair-share	Mature multi-factor priority with configurable decay	Kueue provides basic fair-share with ClusterQueue borrowing
Container overhead	None from a container runtime when jobs run directly on nodes	Configuration-dependent networking overhead and container startup time
Service management	No native support for long-running services	Native Deployment, rolling updates, health checks
Ecosystem	MPI, OpenMP, scientific computing libraries	Cloud-native, microservices mesh, observability stack

The gang-scheduling row of the table compresses a naming collision that is easy to miss. This chapter uses gang scheduling for atomic allocation of a whole execution group, which is the requirement distributed training imposes. Slurm reserves the name for the older, Ousterhout-style meaning of the term, in which co-scheduled jobs share nodes by time-slicing in synchronized rounds. Slurm provides atomic multi-node allocation under its ordinary job model, so the requirement is met even though Slurm's "gang scheduling" option means something else.

The choice of scheduling paradigm also reflects organizational trajectory. Teams starting with HPC-focused training infrastructure may begin with Slurm and add Kubernetes for serving. Teams starting from cloud-native application infrastructure may begin with Kubernetes and add HPC-oriented extensions (Volcano, Kueue) for training. Slurm has meanwhile added a REST interface and container support, so the two platforms now overlap in capability. Either path can reach a functional hybrid architecture, but the starting point determines which capabilities are strongest and which require the most investment to develop.

8.5.5 Hybrid architectures in practice

While the choice between Slurm and Kubernetes often appears binary, engineering organizations can deploy **hybrid architectures** that use the strengths of both systems. The relevant boundary again follows workload semantics rather than labels. Rigid, topology-sensitive jobs fit a batch scheduler's atomic allocation contract, while malleable services fit a reconciliation loop that adjusts replica count continuously. Some training jobs are elastic, and some serving replicas are rigid multi-accelerator gangs, so neither "training" nor "serving" selects a control plane by itself.

The **Kubernetes-over-Slurm** pattern wraps the HPC scheduler within a cloud-native control plane. Kubernetes manages the lifecycle of training jobs through custom operators while delegating batch admission and placement to Slurm. The central engineering problem is ownership: exactly one control plane must authorize each placement, while the outer controller reconciles lifecycle state without issuing a conflicting allocation. The arrangement combines a declarative management surface with Slurm's batch and topology policies, but it also creates a cross-control-plane failure boundary that operators must fence explicitly.

The **Slurm-alongside-Kubernetes** pattern takes a coarser-grained approach, maintaining separate clusters for training and serving that draw from a shared pool of physical nodes. A meta-scheduler or capacity broker acts as an arbitrator, dynamically reassigning nodes between the Slurm partition and the Kubernetes cluster based on demand signals. During intensive training campaigns, nodes are drained from Kubernetes and joined to Slurm; during product launches or traffic spikes, the flow reverses. The primary challenge is the mode-switching cost: draining a node, changing its ownership and configuration, and admitting it to the other control plane takes much longer than an ordinary placement decision. This latency limits sharing to planned capacity shifts rather than second-by-second inference bursts.

A shared application runtime such as Ray runs both training and serving through two abstractions, tasks (stateless function calls executed on remote workers) and actors (long-lived processes that keep state between calls). Ray's placement groups, unlike the cloud network-locality abstraction of the same name, reserve bundles of logical resources for work that must run together, and serving replicas can also require multi-device groups. Inside a shared runtime the workloads keep their different state, latency, and recovery requirements, so the platform must still coordinate Ray's logical allocations with the underlying resource manager and validate the isolation and recovery behavior needed by each workload.

Threading these patterns together, consider the lifecycle of our 175B parameter model. A hybrid platform can allocate the training gang through Slurm and later place serving replicas through Kubernetes. The orchestration boundary concerns resource ownership: the training job releases its allocation, and the serving control plane admits replicas against its own capacity and readiness requirements. Deciding whether the trained artifact is ready for production requires validation and release policy beyond the scheduler; section 12.5 develops those validation checks and rollout policies.

A hybrid platform is useful when its separate allocation contracts justify the additional ownership and recovery boundary. The cost to measure is the coordination between control planes, including node draining and reassignment, rather than a presumed performance advantage from either

product name. Because Slurm and Kubernetes now overlap in capability, the paradigm choice is a workload-specific design decision, and teams must establish the allocation and recovery guarantees they need before layering physical placement constraints on top of them.

Checkpoint 8.1: Scheduling paradigm trade-offs

Before the chapter turns to topology-aware placement, review the core trade-offs:

- ☐ **Gang scheduling**: explain why a rigid synchronous execution group needs atomic allocation, and why the scheduler must still reserve capacity for a waiting gang that holds no GPUs itself.
- ☐ **Slurm vs. Kubernetes**: compare authoritative allocation with desired-state reconciliation, and explain why either design must prevent duplicate ownership during a network partition.
- ☐ **Group admission**: explain why a default scheduler that places pods one at a time can strand capacity for a 64-pod training job, and how a PodGroup's minimum membership prevents it.
- ☐ **Resource units**: explain how the resource unit a device plugin exposes changes what the scheduler can pack, and why a small model on an exclusively assigned GPU strands capacity.
- ☐ **Backfill incentives**: explain why backfill scheduling depends on accurate job runtime estimates and why users have game-theoretic incentives to provide inaccurate estimates.

A control plane can now represent, admit, and reconcile an allocation, yet it still counts accelerators as if any group of the requested size were interchangeable. Two allocations of the same size can meet very different communication budgets depending on where the devices sit in the cluster fabric. The next layer of the scheduler must therefore weigh physical location alongside resource count.

8.6 Topology-Aware Scheduling

A scattered allocation can place frequent collective traffic on shared cross-rack paths, exposing latency or bandwidth limits that a more compact allocation avoids. The penalty is set by how much of the job's traffic crosses those paths and how much capacity the fabric provisions there, not by the rack count itself. The earlier feasibility rules already required compatible topology; topology-aware scheduling makes that requirement explicit in the placement objective, accounting for communication paths as well as raw compute.

8.6.1 The topology hierarchy

Splitting a tensor-parallel group across node boundaries moves its per-layer activation exchanges off intra-node NVLink onto the external network fabric, where an H100's per-direction NVLink rate is 9× the rate of one NDR port (section 2.5.1). Each level of a cluster's communication hierarchy comes from the hardware that builds it and sets a distinct bandwidth and latency budget. GPUs inside a node exchange data over NVLink, the tier fast enough for the per-layer exchanges of tensor parallelism. Nodes under one leaf switch communicate over an external fabric such as InfiniBand or RoCE. Racks joined through spine switches add switch hops whose delivered bandwidth depends on concurrent flows and oversubscription even in a fat-tree. Cross-rack communication also adds tail latency variability: while median latency may be only slightly higher, p99 latency can spike during congestion events, creating straggler effects in synchronous training where the slowest communicator determines the pace for all workers.

A 64-GPU job within one rack can avoid cross-rack paths that a spread allocation must use. It outperforms the spread allocation only when those paths add exposed latency or bandwidth contention; a non-blocking, uncongested spine need not reduce sustained throughput. The performance difference stems directly from the hierarchical communication analysis in section 5.7.1: a flat ring AllReduce runs at the speed of the slowest link in the ring, and cross-rack hops introduce both higher latency and increased contention risk. For synchronous training, every iteration waits for the

↩ Bisection bandwidth and fat-tree network topologies are established in section 3.5.2.

slowest worker to complete its AllReduce, so even occasional cross-rack congestion inflates the Fleet Law's communication term $T_{\text{comm}}(N)$ (equation 1.3) on every step it touches.

The scheduler must therefore model the cluster as a hierarchical graph with discrete bandwidth boundaries (figure 8.7). In the illustrated H100 fleet, per-direction bandwidth steps down from 450 GB/s across the NVSwitch inside a node to 50 GB/s per port under a leaf switch, and to roughly 12.5–25 GB/s per endpoint when balanced traffic saturates an oversubscribed spine. A node boundary can introduce a large bandwidth reduction, but a non-blocking rack or pod boundary need not reduce per-endpoint bandwidth. Oversubscription adds a separate constraint when traffic saturates shared uplinks. Topology-aware placement minimizes the communication cost of the job's actual traffic pattern rather than applying a fixed penalty to every boundary.

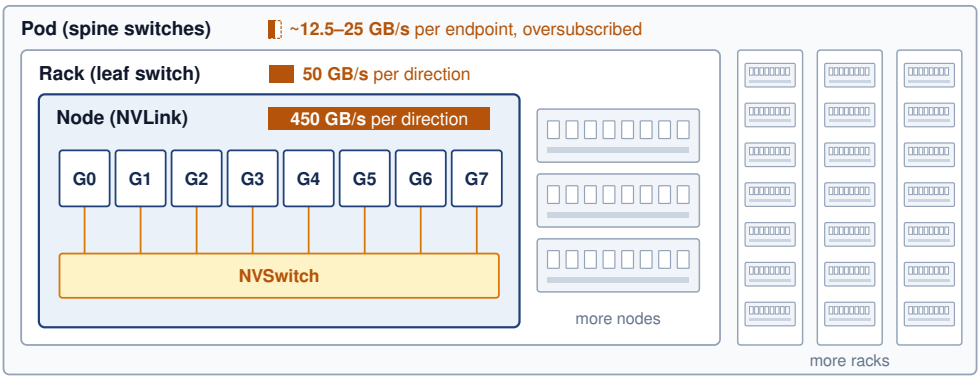

Figure 8.7: Physical Topology Hierarchy and Bandwidth Cliffs: Hardware boundaries dictate communication performance across three distinct domains. The example compares a 450 GB/s per-direction NVLink endpoint budget with a 50 GB/s per-direction external port. Under balanced load saturating a 2:1 to 4:1 oversubscribed spine, the external per-endpoint share falls to approximately 12.5–25 GB/s. A non-blocking spine need not impose this reduction. The bar in each domain's header draws its bandwidth on one shared linear scale.

For a 4-GPU collective group, the routing paths make that cost concrete (figure 8.8). In the illustrated contended fabric, scattering ranks across Rack A and Rack B forces a 3-hop leaf-spine-leaf path with about 12.5 GB/s of effective bandwidth per endpoint, a share of contended uplinks rather than a physical link rate. Consolidating all four ranks (G0–G3) under Rack A's leaf switch avoids those shared uplinks and provides about 50 GB/s per GPU under the example's load assumptions.

8.6.2 Placement algorithms

The illustrated 4× bandwidth difference between the 1-hop leaf path and the 3-hop spine route (figure 8.8) stretches the transfer component by the same factor when transfer time is bandwidth-dominated. Crossing a spine does not impose that ratio independently of load and topology. Whether communication exceeds computation also depends on the model and overlap schedule, and repeated exchanges can amplify the exposed cost. Topology-aware placement algorithms assign jobs to nodes that reduce this cost, turning resource packing into graph optimization over the physical topology. A simple **locality cost**, defined in equation 8.2, penalizes allocations spanning multiple topology domains:

$$\text{Cost}_{\text{locality}} = \sum_{i<j} w(d(g_i, g_j)) \tag{8.2}$$

where g_i and g_j are GPUs assigned to the job, $d(g_i, g_j)$ is their topological distance (0 for same node, 1 for same rack, 2 for different racks), and $w(\cdot)$ is a weighting function that increases with distance. The scheduler minimizes this score when selecting an allocation from the set of feasible placements. An illustrative heuristic assigns locality penalties of $w(0) = 0$ (same node, NVLink, no penalty), $w(1) = 1$ (same rack, one switch hop), and $w(2) = 10$ (different racks, multiple switch hops). These dimensionless weights are policy assumptions rather than measured latency or bandwidth

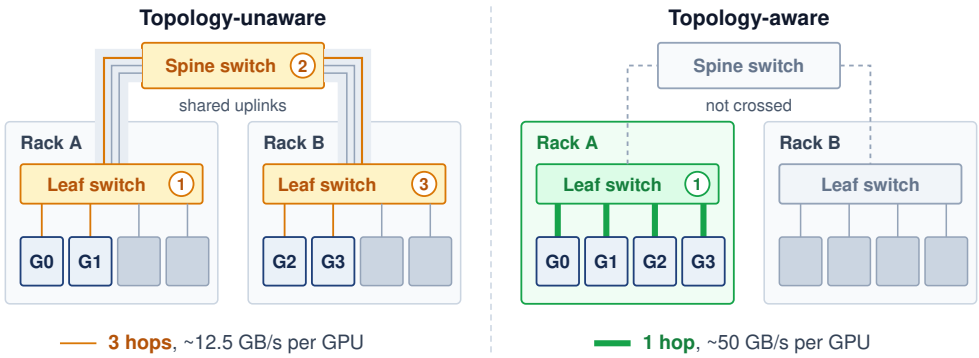

Figure 8.8: Topology-Unaware vs. Topology-Aware Placement: An illustrative 4-GPU job that permits inter-node placement and uses an oversubscribed cross-rack fabric, placed across two racks (left, topology-unaware) must traverse leaf-spine-leaf hops: 3 switch hops per AllReduce message at about 12.5 GB/s effective bandwidth. The same 4-GPU job packed within one rack (right, topology-aware) uses a single leaf switch: 1 hop at about 50 GB/s, roughly 4× faster. Blue squares are the job's GPUs; gray squares are GPUs held by other jobs.

ratios, so a deployment must calibrate them to its traffic and topology. An unweighted pairwise score still cannot distinguish heavily communicating rank pairs from pairs that rarely exchange data.

More sophisticated algorithms distinguish the communication requirements of different parallelism strategies. Frequent tensor-parallel activation exchanges favor a high-bandwidth, low-latency domain that meets the measured communication budget; that domain need not coincide with one server. Pipeline parallelism exchanges activations across stage boundaries rather than between every layer, but its tolerance for link delay depends on microbatch size and the pipeline schedule. Data-parallel gradient synchronization can overlap computation, making wider placement feasible when the exposed communication remains within budget. These patterns guide placement without establishing a universal ranking for every model or fabric.

Algorithm 8.1 combines the all-or-nothing gang constraint of section 8.4.3, the locality cost of equation 8.2, and an illustrative parallelism-to-topology mapping. The algorithm constructs feasible rank mappings, scores each allocation, and chooses the lowest illustrative locality score, leaving the job queued when none fits. This surrogate does not guarantee minimum communication time or completion time. A fleet-scale implementation must restrict that search using heuristics rather than enumerate every allocation.

Algorithm 8.1 Topology-aware placement

Require: job requesting N accelerators with tensor-, pipeline-, and data-parallel groups; cluster topology with illustrative per-level locality weights $w(\cdot)$
Ensure: an all-or-nothing feasible placement with the lowest illustrative locality score, or the job stays queued
1: enumerate candidate slots: sets of N free accelerators that satisfy every resource dimension ▷ *gang: all or nothing*
2: retain only slots that admit a rank mapping satisfying the job's tensor-, pipeline-, and data-parallel communication constraints
3: **if** no feasible slot exists **then**
4: leave the job queued without holding any accelerators ▷ *prevents hold-and-wait deadlock*
5: **return** no placement
6: **end if**
7: **for** each feasible slot **do**
8: use a feasible rank mapping: tensor-parallel ranks within an NVLink domain, pipeline stages within a rack, data-parallel replicas across racks where permitted
9: score the slot by $\text{Cost}_{\text{locality}} = \sum_{i<j} w(d(g_i, g_j))$
10: **end for**
11: **return** the feasible slot of least $\text{Cost}_{\text{locality}}$ ▷ *minimum score among enumerated feasible placements; atomic allocation*

Applied to a whole job, the same reasoning ranks placement strategies. The latencies in this comparison are illustrative, not derived. They are constructed to stay above the ring-AllReduce bandwidth bound for the payload, and they show the order in which placement removes delay rather than values a reader could recompute from topology names. For a 0.5 GB gradient synchronization on 64 GPUs, random placement spans multiple racks and spine switches and takes an assumed 120 ms. Each successive strategy removes a source of delay that the previous one leaves. Confining the group to a single rack removes spine hops and reduces latency to 85 ms. Rail-optimized alignment maps ranks so that every member of each collective group sits at the same local GPU index and therefore on the same network rail (developed in section 8.6.3). Traffic that previously hopped between rails through the spine now stays within one rail's switches, which cuts latency to 40 ms. Topology-optimal placement additionally avoids leaf ports that other jobs are using and reaches 25 ms, a 4.8× reduction (figure 8.9, where each bar also carries its algorithm throughput). This ratio compares whole collectives under four assumed placements, so it need not match the link-bandwidth ratio of the two-rack paths above.

A **hierarchical placement** algorithm applies these constraints from the most restrictive communication domain outward, matching each 3D parallelism dimension to the interconnect tier its traffic can tolerate (figure 8.10). In the illustrated fleet, it places each tensor-parallel group within an NVLink-connected node, where the per-layer activation exchanges run at 450 GB/s. It places consecutive pipeline stages under the same leaf switch, where stage boundaries cross at 50 GB/s. It distributes data-parallel replicas across racks only when the once-per-step gradient synchronization fits the roughly 12.5–25 GB/s per endpoint that an oversubscribed spine leaves. This alignment reduces costly boundary crossings, but placement alone cannot reserve the fabric's full bisection bandwidth against other jobs.

Evaluating all possible placements for a 256-GPU job across a 10,000-GPU cluster is combinatorially intractable, so production schedulers rely on heuristic approaches. Common choices include greedy algorithms that build placements bottom-up (first fill NVLink domains, then fill racks, then spread across racks), tree-search algorithms with pruning that explore the most promising placements first, and scoring-based approaches that evaluate a sample of feasible placements and select the best. Long jobs can amortize a more expensive search, but frequent submissions, short jobs, and recovery bursts also consume scheduler capacity, making end-to-end completion time, including placement delay, the useful comparison rather than execution throughput alone. The same comparison governs how long a job should wait for a compact placement. Holding a job until enough GPUs free up

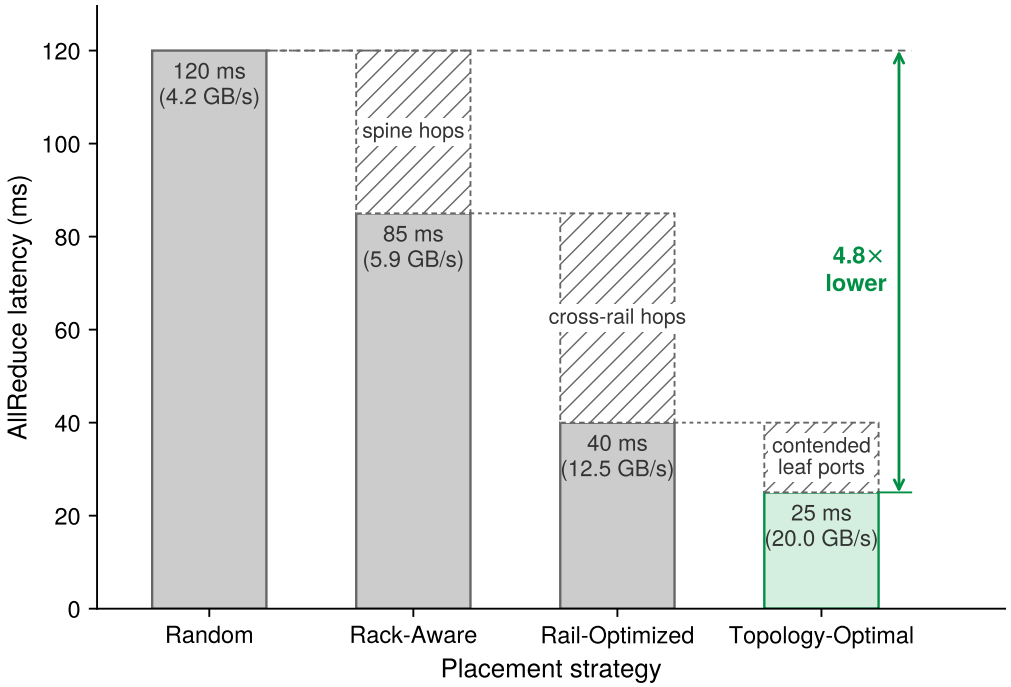

Figure 8.9: Topology Placement Impact: Assumed AllReduce latency across four placement strategies (Random, Rack-Aware, Rail-Optimized, and Topology-Optimal) for a 0.5 GB gradient sync on 64 GPUs. The assumed latency falls from 120 ms to 25 ms, a 4.8× ratio. Labels report payload divided by completion time (algorithm throughput), not physical link bandwidth. The dashed ghost above each bar names the source of delay that strategy removes.

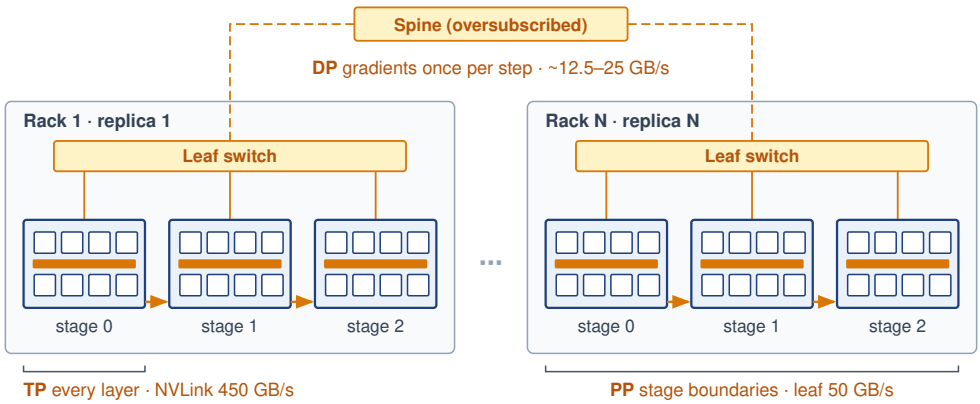

Figure 8.10: Hierarchical Placement for 3D Parallelism: This placement example maps distributed training dimensions to physical interconnect layers. Tensor parallelism (TP), which exchanges activations at every layer, stays inside a node on NVLink (450 GB/s); pipeline parallelism (PP) maps across adjacent nodes under one leaf switch (50 GB/s); and data parallelism (DP), which exchanges gradients only once per step, spans racks through the oversubscribed spine. Bandwidths are per direction, as in figure 8.7.

under one leaf switch delays its start, and the wait pays off only if faster collectives recover more than that delay over the job's run.

Napkin Math 8.2: Placement impact on 3D-parallel training

Scenario: Consider a 256-GPU training job using 3D parallelism: 8-way tensor parallel, 4-way pipeline parallel, 8-way data parallel.

Good placement (topology-aware):

- Tensor parallel groups: 32 groups of 8 GPUs, each confined to one 8-GPU node (NVLink: 450 GB/s per direction)
- Pipeline stages: 4 consecutive nodes within the same rack (1 switch hop, InfiniBand: 50 GB/s per direction)
- Data parallel replicas: 8 racks (3 switch hops through leaf-spine-leaf paths)

Bad placement (topology-unaware):

- Tensor parallel groups split across 2 nodes (InfiniBand instead of NVLink)
- Pipeline stages scattered across racks (3 switch hops instead of 1)
- Data parallel replicas randomly distributed

Math: The tensor parallel communication alone illustrates the impact. With NVLink delivering 450 GB/s per direction, a 1 GB activation transfer takes approximately 2.2 ms. Over InfiniBand at 50 GB/s, the same transfer takes approximately 20 ms, a 9× slowdown.

Systems insight: This calculation prices only the tensor-parallel component of the bad placement. Splitting a tensor-parallel group across nodes makes each activation exchange 9× slower, and the group performs such exchanges at every transformer layer of every step, so the penalty multiplies with model depth rather than appearing once. How much of it reaches step time depends on how much of the exchange was hidden behind computation, and it is why hierarchical placement gives tensor-parallel groups the fastest tier first. The scattered pipeline stages and data-parallel replicas add their own penalties on top, and the scheduler must profile the complete job before assigning a throughput benefit to the placement.

8.6.3 Rail-optimized scheduling

Even when two nodes sit in the same pod, placing corresponding data-parallel shards on mismatched local GPU indices forces their gradient collective to cross between network rails through the spine instead of staying on a single rail switch. In rail-optimized topologies, corresponding GPU positions across nodes connect through parallel network "rails" that a suitable rank mapping can use concurrently. Balanced bandwidth depends on distributing traffic across those paths rather than concentrating it on a subset. Keeping a collective group on one rail avoids cross-rail transfers for that group, while the placement of other groups determines how much rail capacity it must share.

↩ Rail-optimized fabric architectures and rail-to-rail traffic isolation are established in section 3.5.3.

Rail-aware scheduling assigns corresponding model shards across data-parallel replicas to GPUs attached to the same rail. When intra-node parallelism partitions a model replica across the eight GPUs of a node, each data-parallel gradient collective reduces only one shard i across replicas rather than combining all eight GPUs on the node. In the eight-GPU-per-node example, local GPU index i connects to rail i. Placing every replica's copy of shard i at that same local index on its selected node means the data-parallel replica rank varies across nodes rather than across local GPU indices, so the collective for shard i stays entirely within rail i. This alignment removes cross-rail transfers for that group, though other groups or tenants can still contend for the same rail's links, and full dedicated bandwidth requires an explicit capacity and sharing contract.

Rail alignment therefore acts at two levels, and they belong to different components. Within a node it is a rank-mapping decision. The launcher or framework must give every member of a data-parallel group the same local GPU index on its node, which is what removes the cross-rail hops in the placement comparison above. A mapping that puts shard i at index i on one node and at another index elsewhere sends that group's traffic between rails.

Across nodes, rail alignment constrains the scheduler only when allocations are smaller than a node. When each replica's pipeline stage occupies a complete node and every node wires local GPU i to rail i, any free node already offers every rail, so rail alignment does not narrow which racks

the scheduler can choose. With partial-node allocations, a job can receive different local indices on different nodes, and the scheduler must then choose GPU positions as well as nodes so that each group's members share a rail. Those positions can conflict with the tensor-parallel budget and the paths between pipeline stages when only particular GPU positions remain free. The packing cost therefore follows the allocation granularity and the available topology rather than an inherent conflict among the three parallelism dimensions. Even a single-node allocation can have meaningful topology constraints when GPUs occupy different interconnect, PCIe, or non-uniform memory access (NUMA) domains, so packing efficiency should be optimized only among placements that satisfy the job's local and external communication contract.

8.6.4 Topology discovery and dynamic adaptation

When a spine switch fails or an optical link degrades and begins retransmitting, a static configuration file such as Slurm's `topology.conf` or Kubernetes node labels still reports the rack as healthy and full-bandwidth. The scheduler then places communication-heavy gangs onto a crippled path. Runtime **topology discovery** supplements static maps with information exposed by the operating system and device libraries. NCCL uses system topology information, including the PCIe hierarchy and network devices exposed through paths such as `/sys/class/infiniband`, to identify communication paths and network interface card (NIC) proximity. The discovered graph describes reported connectivity; health checks and traffic measurements must establish whether those paths deliver the expected performance.

Cluster topology is an operating state, not a fixed schematic. Hardware failures, maintenance windows, and thermal throttling change the paths and effective capacity available to the scheduler. A spine switch failure in a redundant Clos network can leave connectivity intact while removing capacity. The effect on our 175B parameter model's data-parallel AllReduce depends on rerouting, concurrent traffic, and how much bandwidth the job was using. If the failure increases exposed communication time, the scheduler must compare continuing at reduced throughput with the cost of moving to a healthier allocation. A degraded link can also pass binary up/down probes while cutting delivered bandwidth, so the scheduler needs performance evidence such as retransmission counts and collective bandwidth, the link-level telemetry of section 3.9.1, as well as reachability.

A checkpoint-and-restart migration spends time saving state, releasing the old allocation, obtaining a new one, and restoring execution. Its value depends on measured transition cost and remaining work. If migration takes 10 minutes and removes a sustained 20 percent throughput loss, each later minute of training recovers a fifth of a minute of progress, so the move pays back in less than an hour of subsequent training. That comparison assumes a healthy destination is available and the gain persists. A practical policy also accounts for queue delay and requires evidence that placement causes the slowdown, such as rising collective time on the affected path, before it moves the job.

War Story 8.2: Performance-aware health checks (2021)

Context: OpenAI's account of scaling Kubernetes describes automated health checks that remove misbehaving nodes from the schedulable pool (Sigler 2021).

Mechanism: Passive checks monitor reachability, disks, and GPU errors, but error codes do not expose every hardware or driver problem. Active GPU tests exercise the device before admitting ordinary workloads and periodically during its lifetime.

Fix: Nodes enter with a preflight taint, a node marking that repels ordinary pods, which blocks normal scheduling until the tests pass. Failed passive checks cordon a node, marking it unschedulable for new work; more serious failures can trigger eviction requests.

Systems lesson: Reachability alone is not a sufficient admission test. A scheduler needs health evidence that reflects the resources its jobs depend on, together with a policy for removing suspect nodes from placement.

The general scheduling idea is quarantine: when runtime evidence says a node, rack, or link is degraded, the scheduler should stop placing sensitive work there even if the nominal resource count still looks available. Kubernetes implements this pattern through taints and tolerations,[9] while

[9] **Taints and tolerations**: Node taints can block scheduling of pods without matching tolerations. A toleration permits consideration of a node but neither requires placement there nor reserves it exclusively for training. Dedicated GPU pools also need controlled toleration assignment and placement constraints; removing existing workloads requires an eviction or drain policy appropriate to the taint's effect.

Slurm clusters commonly use drain states, node features, or partition rules.

↩ Slow components that pass liveness checks while setting the pace of every synchronous step, and how to tell them from fail-stop failures, are established in section 7.9.5.

The complementary problem is how jobs adapt when available resources change dynamically. Rather than treating resource allocation as fixed for a job's lifetime, elastic training allows logical jobs to grow, shrink, and recover.

8.7 Elastic Scheduling

When cluster demand fluctuates, rigid allocation creates two categories of waste. During off-peak hours, a job running on 256 GPUs cannot absorb idle resources to speed up training. During peak demand, the same job cannot release excess resources to accommodate higher-priority work without being fully preempted. The scheduler's only options are "let it keep all its resources" or "kill it entirely," with nothing in between.

Elastic scheduling expands this binary choice. The recovery mechanics of elastic training show how a distributed job can adjust its worker count, adapt its batch policy, and reconstruct its communication group after a failure. From the scheduler's perspective, the same capability solves a resource-allocation problem under fluctuating demand. A training job might start with 128 GPUs, scale up to 256 when resources become available, and scale back to 192 when higher-priority work arrives, provided each allocation supports a viable worker group. Resizing may pause training and restart workers from saved state. Whether worker processes restart depends on the framework; TorchElastic restarts its worker group when membership changes. The scheduler therefore chooses among supported allocation sizes, each with a transition cost, rather than assuming uninterrupted progress at any resource level.

↩ Batch-size and learning-rate adjustment and communication-group contraction during elastic recovery are established in section 7.10.1.

8.7.1 The scheduler-framework contract

If a scheduler reclaims four nodes from a running job by abruptly killing their containers, or adds four idle nodes without triggering a new rendezvous, the surviving ranks treat the change as a fatal collective timeout rather than an elastic resize. Resizing a live job without crashing it requires an explicit contract between the scheduler and the training framework. For planned resizing, the scheduler and framework coordinate resource release through defined signals and an agreed adaptation window. Unexpected failures can bypass that window, so the framework must also recover from valid saved state. Its recovery path reconstructs communication groups and either preserves the global batch or applies a validated batch-size and learning-rate policy. The worker-membership lifecycle therefore behaves as a state machine centered on active training (figure 8.11). Idle GPUs trigger scale-up, preemptions trigger scale-down, and host timeouts or ECC errors divert ranks to worker replacement; each path returns to active training once the communication group is rebuilt or a spare node joins.

From the scheduler's perspective, this contract is defined by three governing parameters: the elastic range, the transition cost, and the scaling efficiency curve. The elastic range, $[N_{\text{min}}, N_{\text{max}}]$, specifies lower and upper worker bounds; the job's parallelism layout determines the valid allocation increments within them. The transition cost includes any checkpoint, restart, state redistribution, and rendezvous required by the runtime. Measuring the full interruption determines how frequently the scheduler can trigger reallocations before transition overhead wipes out any throughput advantage. Finally, the scaling efficiency curve quantifies throughput at each worker count relative to linear scaling. The marginal throughput of each added worker follows from that curve and indicates whether assigning idle GPUs to this job yields more aggregate progress than holding them for an assembling gang-scheduled workload. Together, these parameters provide the scheduler with the quantitative bounds needed to resize jobs dynamically without treating elasticity as free capacity.

The scheduler must also account for the effect of worker count on global batch size ($B_{\text{global}} = B_{\text{per-worker}} \times N$), where the per-worker batch counts examples contributing to one optimizer update, including gradient accumulation. Holding the global batch constant requires each remaining worker to process more examples after a shrink. Larger microbatches consume more memory; additional accumulation steps extend the update time. Preserving the batch removes one source of optimization change but does not guarantee identical convergence. Allowing the global batch to change instead requires validating the learning-rate policy and resulting model quality, so throughput alone cannot determine the best allocation.

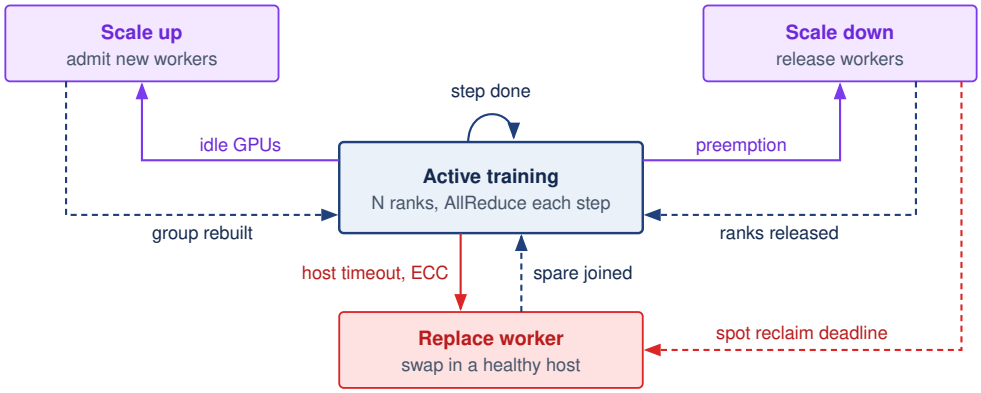

Figure 8.11: Elastic Training Worker Lifecycle and State Machine: The state machine governing dynamic worker membership. Active training loops on step completion and leaves that state along three paths: idle GPUs trigger scale up, a preemption for higher-priority work triggers scale down, and a host timeout or ECC error triggers worker replacement. Each path rebuilds the communication group and returns to active training (dashed arrows), and a spot reclamation deadline can force a shrinking job straight into replacement. Worker processes may restart during resizing, so recovery requires the application to restore valid training state.

Listing 8.2: TorchElastic Launch Configuration: The `--nnodes=4:16` range configures the elastic worker group to admit 4 to 16 nodes, or 32 to 128 workers. A separate cluster scheduler or operator must acquire, launch, and release nodes within that range.

```
# Launch elastic training with min 32 to max 128 workers
torchrun \
  --nnodes=4:16 \
  --nproc_per_node=8 \
  --rdzv_backend=c10d \
  --rdzv_endpoint=head-node:29500 \
  --max_restarts=3 \
  train.py
```

TorchElastic sets the elastic range in its launch configuration as a minimum and maximum node count (listing 8.2). There, the `--nnodes=4:16` specification configures TorchElastic's rendezvous and worker-membership range as 4 to 16 nodes (32 to 128 workers with 8 processes per node). It does not acquire or resize cluster resources; a separate cluster scheduler or operator must place the nodes and start or stop elastic agents as capacity changes.

Systems Perspective 8.2: Elastic vs. rigid scheduling

Consider a congested cluster where a 128-GPU data-parallel job, with no model-parallel constraint on its worker count, waits 8 hours for a full gang allocation. This first comparison assumes throughput scales linearly with worker count, equivalent learning progress per processed example, and negligible resize overhead; the later scaling decision relaxes the throughput and transition assumptions.

Rigid scheduling: Wait 8 hours, then train for 24 hours at full throughput. Total wall-clock: 32 hours.

Elastic scheduling (min 32, max 128):

- Start immediately with 32 GPUs (throughput = 25 percent of full)
- At hour 2, when 64 GPUs free, scale to 64 (throughput = 50 percent)

- At hour 8, when the full gang frees and the rigid job would also start, scale to 128 (throughput = 100 percent)

Work completed during the 8-hour queue wait, in hours of full-scale training:

- Hours 0 to 2: 25 percent throughput $\times$ 2 hours = 0.5
- Hours 2 to 8: 50 percent throughput $\times$ 6 hours = 3
- Total: 3.5 hours of equivalent full-scale work

The elastic job completes in approximately 28.5 hours (8 hours of partial-scale training plus 20.5 hours at full scale), saving 3.5 hours compared to rigid scheduling. This improvement of about 11 percent comes entirely from using resources productively during what would otherwise be idle queue time.

The scheduler cannot treat every worker count as a valid elastic configuration. Model parallelism constrains the allocation increment when the resize policy keeps tensor and pipeline layouts fixed and adds or removes whole data-parallel replicas. For a 175B-parameter job that relies on 8-way tensor parallelism within each node and 8 pipeline stages across nodes, this policy operates at the granularity of a whole model replica of 64 GPUs. The pipeline depth of 8 is the smallest power of two at which each GPU's share of the mixed-precision Adaptive Moment Estimation (Adam) training state fits in its HBM when that state is not sharded across replicas. The layout of chapter 6 instead uses a shallower pipeline and shards optimizer state across the data-parallel group with ZeRO (section 6.3.3). This job gives up that sharded layout on purpose. Keeping each replica's state self-contained costs memory, but it lets the scheduler add or remove a replica without redistributing every other replica's optimizer shards. The 175B examples in section 8.7.3 assume these self-contained replicas. The job can scale between 128 nodes (1,024 GPUs) and 32 nodes (256 GPUs) one replica at a time, changing the data-parallel degree. Removing a single GPU from a tensor-parallel group would instead require a new sharding layout and redistribution of training state. Changing the pipeline degree would also require repartitioning stages and rebuilding the pipeline schedule. A runtime that supports these transitions can expose them to the scheduler, but their feasibility and measured costs belong in the resize contract rather than following automatically from data-parallel elasticity.

The transition cost also constrains the scheduler's resize frequency. Each scaling event can interrupt useful work while the framework reconstructs its communication group and restores or redistributes state. The duration must be measured for the chosen runtime and checkpoint path. The scheduler must amortize this cost over the subsequent run duration, so small resize increments are only worthwhile if the job will run at the new scale for significantly longer than the transition window, a test that section 8.7.3 applies to resize policies.

8.7.2 Scheduler integration

Without scheduler integration, a job that declares an elastic range $[N_{\text{min}}, N_{\text{max}}]$ in its launch script still waits in the queue for its full node count and still gets killed outright when a higher-priority job arrives. Integrating the $[N_{\text{min}}, N_{\text{max}}]$ bounds into the scheduler lets it start a job early at the minimum, expand the job into newly idle capacity, or shrink it to release borrowed capacity without terminating the logical job. Which action wins depends on the job's scaling curve and transition cost, and the policies below compare the choices by the work completed over the same time window.

Each action must respect the framework contract. The most consequential of the three is **graceful preemption**: when higher-priority work arrives, the scheduler first shrinks an elastic job toward N_{min}, reclaiming whole replicas at the next resize boundary, and resorts to checkpoint-and-preempt only when the final viable group itself must yield. Quota borrowing across teams, developed in section 8.11, depends on this primitive whenever an owner reclaims capacity it lent. Scaling curves and transition costs determine whether the extra progress repays the disruption. Because batch changes can affect convergence, the test of a resize is time to target model quality, not throughput. Throughput counts examples per second, but a changed global batch can change how many examples the run needs to reach its target quality, so an expansion helps only when the faster rate outweighs any added examples.

8.7.3 Elastic scaling policies

Adding resources to a running job does not always accelerate time-to-convergence. Communication and synchronization overhead limit throughput gains, requiring the scheduler to enforce elastic scaling policies based on empirical efficiency curves rather than resource availability alone. With fixed global work per step, a job's scaling efficiency $\eta_{\text{scaling}}(k)$ at k workers (definition 6.4) typically passes through three regimes. In the near-linear regime, per-worker computation dwarfs every fixed per-step cost. In the sublinear regime, pipeline fill and drain and the unhidden part of gradient synchronization become a visible share of each step. In the saturation regime, those fixed costs rival the shrinking computation itself. Assigning available GPUs to a saturated job can yield less useful progress than serving queued work. The lower allocation bound answers a different question. Below its **minimum viable scale**, a job may lack the memory or complete parallel groups needed to execute at all. A policy can impose a higher minimum to meet a deadline, even when a smaller allocation remains computationally efficient.

The 175B reference job shows these regimes with the chapter's own numbers. Modeled at a fixed global batch with the step-time decomposition of the Fleet Law (equation 1.3), its scaling efficiency falls from 87 percent at 256 GPUs to 83 percent at 1,024 GPUs and 69 percent at 4,096 GPUs (figure 8.12, panel a). Efficiency falls, yet quadrupling from 1,024 to 4,096 GPUs still raises throughput 3.3×, so a lower efficiency is a declining marginal return rather than a slowdown. Panel (b) locates the loss. Tensor-parallel communication shrinks along with each worker's computation, so its share of the step never rises above 11 percent. The decline comes instead from pipeline fill and drain and unhidden gradient synchronization, which take roughly fixed time per step while each worker's share of the batch shrinks. The pipeline bubble grows fastest of the three, for a reason the bubble fraction $(N_{\text{PP}}-1)/(m+N_{\text{PP}}-1)$ of section 6.5.2 makes explicit. With the global batch fixed, each added replica takes a share of the sequences, so the microbatch count m per replica falls while the stage count stays at 8, and the idle fraction rises.

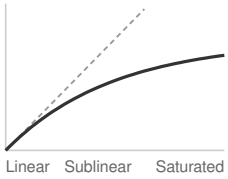

Throughput falls away from ideal linear scaling and saturates.

The shaded window from 256 to 1,024 GPUs is therefore a policy choice, not an interval derived from these efficiencies alone. A minimum allocation could follow from a completion deadline or the job's declared resource contract. It does not follow from insufficient per-worker computation at smaller scale, since with fixed global work fewer workers generally perform more work each. The upper limit likewise requires comparing additional useful progress with resize cost and the value of serving queued jobs. A real scheduler must obtain those inputs from profiling and the workload's quality target.

This logic extends to the temporal dimension, where the scheduler must weigh the immediate progress of a smaller allocation against waiting for a larger one. If a 512-GPU job can start immediately with 128 GPUs or wait for the full request, the better choice depends on the integration of throughput over time. The resize overhead (the cost of stopping, checkpointing, and restarting when the runtime requires them) must be amortized over the subsequent run duration. A policy of "opportunistic scaling" that resizes a job every time a single node becomes free can lose its throughput advantage to repeated interruptions, completing less useful work than a stable allocation.

> Napkin Math 8.3: The elastic scaling decision
>
> The comparison of elastic and rigid scheduling above assumed linear scaling and free resizing; this decision drops both assumptions. Consider a training job with a target batch size that requires 512 GPUs for the highest throughput among the configurations considered. Three scheduling strategies are evaluated over a window of 24 hours.
>
> **Given**: Throughput at 128 GPUs is 1 epoch/hour (baseline). Throughput at 256 GPUs is 1.8 epochs/hour (90 percent scaling efficiency relative to that baseline). Throughput at 512 GPUs is 3.2 epochs/hour (80 percent scaling efficiency). Re-scaling cost is 10 minutes (0.17 hours) to checkpoint, re-shard, and restart.
>
> **Strategy 1 (immediate start, no scaling)**: The job starts immediately with 128 GPUs and runs for the full 24 hours. Total work: $24 \times 1.0 = 24$ epochs.

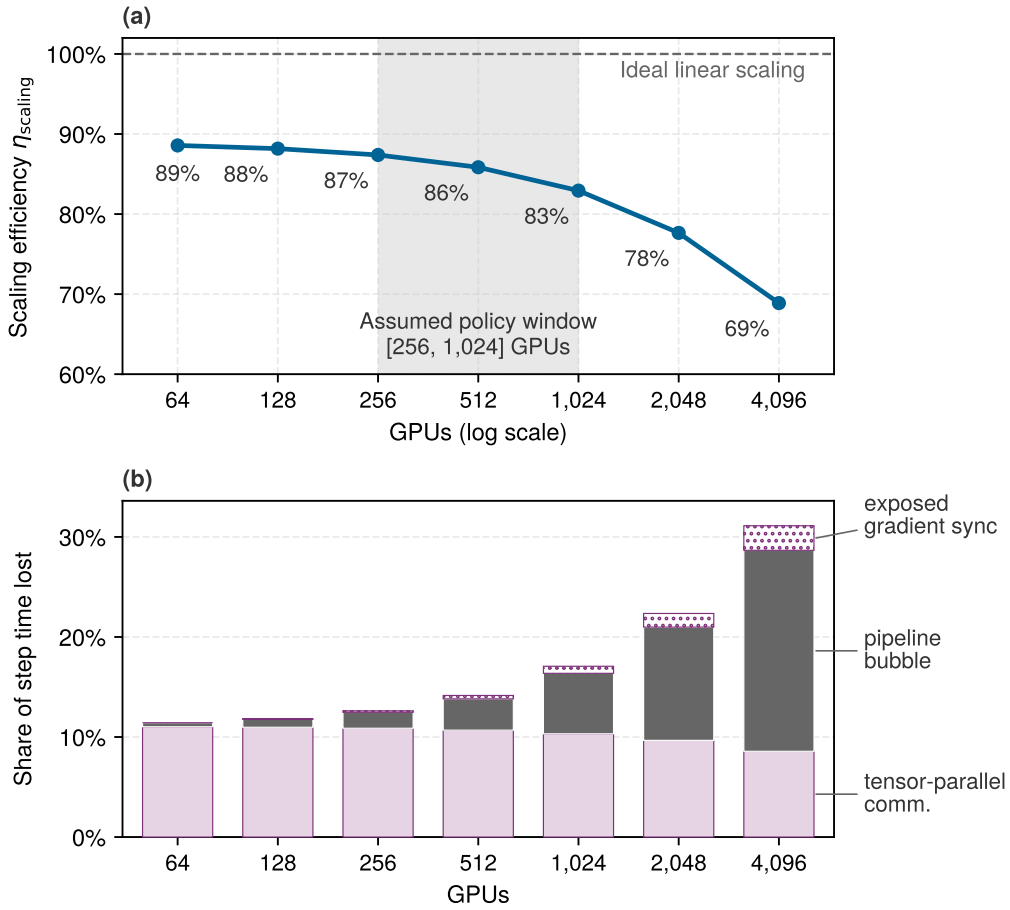

Figure 8.12: Modeled Scaling Efficiency and Elastic Operating Range: (a) Scaling efficiency (η_{scaling}) of the 175B reference job from 64 to 4,096 GPUs, computed with 8-way tensor parallelism within each DGX H100 node, 8 pipeline stages, HDR InfiniBand, and a fixed global batch of 1,536 sequences. The shaded window is an assumed admission policy; efficiency alone establishes neither its minimum nor its maximum. (b) Share of each step lost to tensor-parallel communication, pipeline bubble, and exposed gradient synchronization. Lower efficiency does not imply that adding workers reduces throughput.

Strategy 2 (wait for the full allocation): The job waits in the queue for 4 hours until 512 GPUs become available, then runs for 20 hours. Total work: $20 \times 3.2 = 64$ epochs.

Strategy 3 (elastic step-up): The job starts immediately with 128 GPUs. After 4 hours, the capacity that lets Strategy 2 start frees up, but the elastic policy caps growth at 256 GPUs so that the remaining GPUs stay available to queued jobs. The job pauses, scales up, and resumes. Phase 1: $4 \times 1.0 = 4$ epochs. Phase 2: $(24 - 4 - 0.17) \times 1.8 \approx 35.7$ epochs. Total work: 4 + 35.7 = 39.7 epochs.

The elastic strategy completes nearly 1.7× the epoch-equivalent work of the immediate small-scale run, but still trails the full allocation by about 38 percent. That comparison counts this job's epochs, not the GPU-hours spent on them. Strategy 2 holds twice as many GPUs as Strategy 3 for most of the window and runs them at 80 percent scaling efficiency against Strategy 3's 90 percent. The elastic job therefore does more work per GPU-hour, and the GPUs it leaves free serve queued jobs. Which strategy wins therefore depends on whether the fleet values

this job's completion time or the total work completed across all jobs. That gap shrinks when queue waits grow or peak capacity stays scarce. Re-scaling overhead is only 10 minutes, so the step-up strategy keeps its lead over the immediate small-scale run unless slow checkpoint transfer or compilation stretches that overhead beyond about 9 hours.

Placement and elasticity determine which allocations are feasible and how much useful progress they produce across competing jobs. Fleet economics adds the price of those choices. Every placement, preemption, and reservation consumes capacity that could serve another workload, so the scheduler must compare both useful progress and opportunity cost.

8.8 Cost Optimization

An idle accelerator in an enterprise fleet incurs operational costs every hour whether computing or blocked in a queue, turning placement and preemption decisions into direct financial trade-offs. The orchestrator translates this utilization-to-dollars relationship from section 8.1 into an online scheduling signal, balancing resource prices, reservation commitments, interruption risk, checkpoint cost, and job deadlines to decide whether to admit, delay, preempt, or resize work. Capital planning, portfolio return, cost attribution, and lifecycle total cost of ownership are larger, slower decisions developed in section 12.9; the scheduler consumes their outcomes as prices and constraints.

At a fixed utilization, idle capacity grows linearly with fleet size, so the dollar cost of a scheduling inefficiency grows with every procurement (figure 8.13). At 10,000 GPUs operating at 50 percent utilization, annual waste reaches $87.6M under the chapter's $2/GPU-hour accounting scenario. Published measurements show that such underutilization occurs in production. Microsoft's Philly cluster study (Jeon et al. 2019) measured mean GPU utilization of approximately 52 percent. The 30 to 70 percent shaded band in the figure is an illustrative operating range for the cost model, not a universal empirical claim about all production clusters. Moving from 50 percent to 80 percent utilization, the improvement priced in the chapter opening, recovers over $52M annually in equivalent capacity value. That is the same 3,000 GPU-equivalents the opening counted as the recoverable 30 percent idle share, and it becomes cash savings only when it avoids billed usage or a planned purchase.

Recovering that idle capacity requires pricing each scheduling lever against the disruption it causes. Because reclaiming GPUs from running jobs is the most direct way to admit higher-value work, the cost of a single preemption and the cascades it can trigger come first. Discounted spot capacity then lowers the hourly price of restart-tolerant training, diversification and checkpoint policies keep correlated reclamations from erasing that discount, and capacity-reservation tiers combine with workload cost policies to route jobs across reserved, on-demand, and interruptible pools.

8.8.1 Preemption cost and cascades

When a high-priority job enters a saturated cluster, the scheduler must decide which running workloads to terminate to free up resources. This decision is rarely isolated. In tightly packed clusters, evicting a medium-priority job to accommodate a high-priority request often triggers a **preemption cascade**, where the evicted job immediately attempts to reschedule itself by displacing lower-priority workloads. The disruption multiplies with each level. If the urgent job displaces k medium-priority jobs and each of those, on rescheduling, displaces k lower-priority jobs, one admission interrupts $k + k^2$ jobs after two levels. Every one of them pays the preemption overhead quantified below and adds a checkpoint reload to the storage system at the same moment. Two dampening controls cap that growth. A rule that evicted jobs requeue without preempting anyone themselves stops the chain at the first level, so one admission costs at most k interruptions. A **preemption budget** limits how often any single job, or what fraction of capacity, may be displaced in a time window, which protects the jobs at the bottom of the priority order that every chain would otherwise reach.

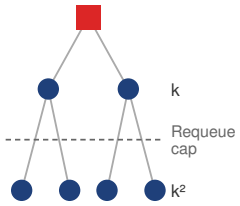

Each eviction can evict again, so one admission reaches $k + k^2$ jobs.

The engineering cost of preemption extends beyond scheduling latency. The **preemption overhead** depends on the exit path. A successful graceful checkpoint adds the time to save and restart while preserving completed progress; a forced stop can instead discard work since the last durable

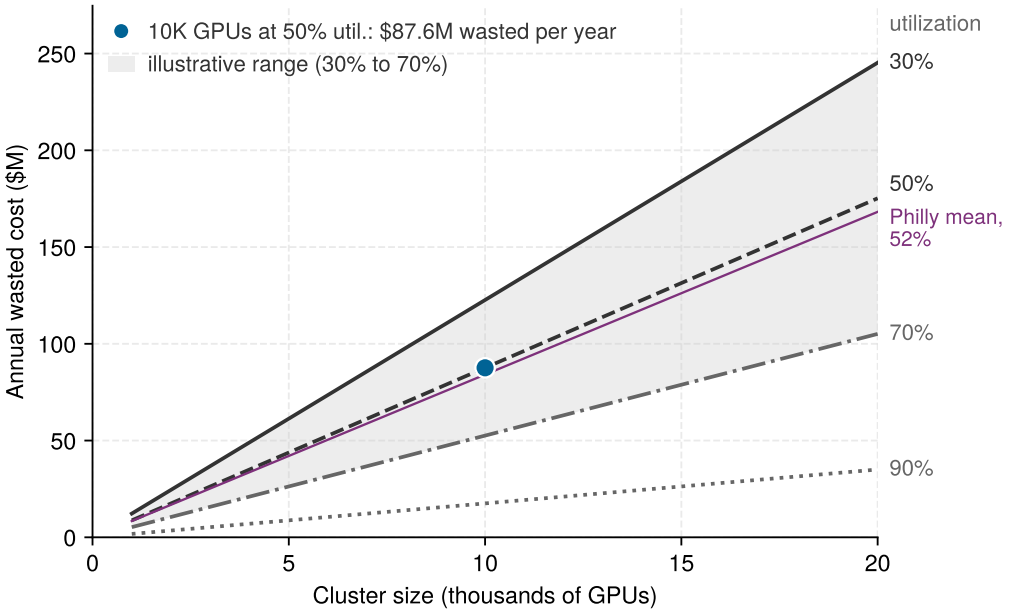

Figure 8.13: The Cost of Idle GPUs: Annual wasted cost as a function of cluster size for different GPU utilization levels, assuming $2/GPU-hour (cloud equivalent). The shaded region between 30 percent and 70 percent is an illustrative operating range for the scenario analysis. At 10,000 GPUs with 50 percent utilization, the model yields $87.6M in annual wasted capacity. The thin violet line traces the published Microsoft Philly mean GPU utilization of 52 percent across all cluster sizes; it is a utilization reference, not a measured cluster size or cost. Data source: Jeon et al. (2019).

checkpoint. Consider such a victim, a forced interruption of the 64-GPU fine-tuning run for our 175B parameter model. The scenario assumes lost work equivalent to 15 minutes of execution. The restart then repeats the cold-restart steps of the recovery budget in section 7.9.2, without its detection step because the scheduler itself stopped the job. These steps take 3 minutes to relaunch, 24.5 seconds to reload the 2.45 TB checkpoint with its optimizer state, and 2 minutes of warmup as data pipelines refill and just-in-time (JIT) compilation and runtime initialization complete. The calculation prices replay and restart across the same allocation; it does not also charge for a successful final save.

Reconstructing the distributed data pipeline also belongs in the preemption contract. The checkpoint must preserve the sampling and loader state needed to resume the intended data order, alongside model and optimizer state. If the restart path omits that state, it can repeat or skip samples and lose exact replay of the training trajectory. A different shuffle order does not by itself prove statistical corruption, but the recovery policy must specify which changes are acceptable. The scheduler should therefore measure restoration and pipeline warmup as part of recovery cost, rather than assume that reloading model weights alone makes the job productive.

The resulting 20.4-minute overhead combines discarded execution, its larger part, with relaunch, reload, and warmup time; it is not a measurement of continuous kernel activity after eviction. Spread over the run, it belongs to the Fleet Law's coordination term $T_{\text{sync}}(N)$ (equation 1.3), which counts scheduling decisions and failure recovery alongside barrier waits, so frequent preemption lengthens the average step exactly as frequent failure does. At $2/GPU-hour, the scenario values each event at $44 in allocation time. A low-priority run sits at the bottom of every cascade, so without a budget it can be struck many times a day. At the scenario's 12 forced preemptions per day, it loses 4.1 hours of equivalent allocation time, or 261 GPU-hours for this 64-GPU run. A per-job budget bounds that loss directly. When a budget is exhausted, the scheduler must consider waiting, another placement, or a defined emergency exception. The trade-off is reduced recovery churn against delayed admission, measured under the workload's deadlines.

8.8.2 Spot instances and preemptible VMs

Renting on-demand cloud GPUs for a multi-week training run protects the job from provider reclamation, but pays full price even when spare GPUs sit idle in the provider's data center. Renting interruptible capacity cuts the hourly rate by 60 to 90 percent, but lets the provider reclaim those GPUs on best-effort notice. Cloud providers sell this spare capacity through spot instances[10] (AWS) or **preemptible VMs** (GCP), where actual discounts depend on the resource and market. The reference interruption windows range from up to 30 seconds for GCP preemptible shutdown to 2 minutes for AWS stop or termination notices, with best-effort delivery rather than a guaranteed recovery interval. At the scenario rates of \$0.70/GPU-hour instead of \$2/GPU-hour, the worked comparison below shows how savings can exceed half the on-demand cost after accounting for its assumed interruption overhead.

[10] **Spot instances**: The scenario's 60 to 90 percent discount purchases reclaimable capacity, not a guaranteed interval for saving state. AWS's best-effort stop or termination notice targets 2 minutes, while GCP preemptible shutdown allows up to 30 seconds; AWS hibernation begins immediately. If a job's measured checkpoint time exceeds the reference 30-second window, a notice-triggered save cannot be its only recovery path.

The viability of spot instances for ML training depends on two factors that determine whether the discount translates into actual savings. The first is the cost of each interruption (checkpoint overhead plus restart time plus lost work since the last checkpoint). The second is the frequency of interruption (which varies by instance type, region, time of day, and overall cloud demand). Equation 8.3 gives the **effective cost** of spot training:

$$C_{\text{effective}} = C_{\text{spot}} \times \frac{T_{\text{total}}}{T_{\text{productive}}} \tag{8.3}$$

where T_{total} is billed allocation time, including checkpointing, recovery, and work later discarded after an interruption, and $T_{\text{productive}}$ is the equivalent training time represented by retained progress. The ratio assumes a fixed allocation size and price, with unchanged useful throughput while training; variable allocations require summing the cost of each allocation interval. Waiting for replacement capacity extends completion time but contributes to billed time only while resources remain allocated. When checkpointing and interruption losses consume a small fraction of billed time, effective cost stays near the spot price. Frequent interruptions, expensive recovery, or substantial replay can instead erase the discount.

Spot savings and fault tolerance infrastructure are coupled. Fault tolerance mechanisms (frequent asynchronous checkpointing, fast checkpoint loading, elastic recovery) make spot capacity usable by lowering the cost of each interruption. Elastic training can resume a logical job with fewer workers after reclamation, provided the remaining allocation is viable, although the framework may still restart worker processes and restore checkpointed state, so elasticity reduces recovery cost rather than eliminating it. The same infrastructure therefore earns a double return. It protects against hardware losses, and it enables the deliberate use of interruptible capacity. Both benefits require measurement. The scheduler combines observed interruption rates, recovery costs, and retained training progress with resource prices to estimate whether spot savings repay the additional recovery infrastructure. The same evidence supports valuing protection against hardware failures; neither benefit follows from the discount or checkpoint mechanism alone.

8.8.3 Spot instance strategies for ML training

Spot capacity pays off only when an interruption rarely removes the whole job at once, so schedulers that use it diversify where capacity comes from instead of simply retrying on failure. Availability zone diversification reduces the probability of simultaneous fleet-wide reclamation. Cloud providers typically manage spot pools independently per availability zone (AZ). If a training job requires 1,024 GPUs, allocating them as a single block in one zone exposes the job to a complete stop if that specific zone's spot pool is reclaimed. Diversifying independent jobs or restart destinations across zones reduces exposure to one exhausted pool. Splitting a single synchronous gang across zones also changes its communication paths, so the scheduler must include cross-zone bandwidth, latency, and transfer costs before accepting that placement. The scheduler still must model correlation carefully: independent per-instance interruptions make large simultaneous losses unlikely, but pool-level or AZ-level reclamations can remove hundreds of GPUs at once and must be treated as elastic-shrink or failover events rather than ordinary isolated failures.

Instance type diversification expands the set of capacity pools a job can use, though availability across those pools can remain correlated. A job that requests only one instance type cannot use spare capacity in another pool. Accepting alternatives therefore requires a compatibility list that includes

memory capacity, runtime support, communication paths, and measured throughput. An instance with more memory can still be an unsuitable substitute if its runtime support, communication paths, or throughput differ. The scheduler compares each feasible alternative's price and recovery cost per unit of retained progress; mixing devices within a synchronous gang also requires accounting for unequal worker speeds.

Spot capacity signals can guide placement but do not necessarily predict interruption frequency. For example, the AWS Spot Placement Score estimates the current likelihood that a capacity request will succeed in a region or availability zone; it explicitly does not predict interruption risk. A scheduler can use this signal to choose where to request capacity, but estimating interruptions per day requires separate historical measurements for the instance type and region. If the resulting workload-specific recovery model shows that interruption overhead erases the discount, the scheduler should migrate the job to on-demand capacity or a different region.

Checkpoint frequency optimization for spot instances extends the reliability logic. The Young-Daly formula developed in section 7.8.1 models losses from unpredicted interruptions, whether caused by hardware failure or reclamation without a usable warning. Provider notices offer an additional opportunity to save state, using the reference windows of 2 minutes on AWS and up to 30 seconds on GCP preemptible VMs. A **two-tier checkpointing** approach combines periodic checkpoints, with an interval chosen for the workload's interruption model, and a best-effort checkpoint triggered by a notice. The notice-triggered checkpoint reduces lost work only if it becomes durable before reclamation; otherwise, the job restores its last durable state, whether produced by a periodic, asynchronous, or incremental checkpoint, and replays subsequent work. That contingency planning matters more when spot interruptions are correlated across a pool rather than isolated to single machines.

Systems Perspective 8.3: Correlated spot risk

Context: Spot capacity is cheap because it is reclaimable. That risk is not independent across instances: cloud regions, instance types, product launches, and deadline cycles can create correlated demand.

Failure mode: A fleet that concentrates all interruptible training in one region and one instance type can lose a large fraction of workers at once, even if each individual worker appears restartable.

Consequence: Checkpointing handles isolated interruptions; it does not solve fleet-wide preemption when every retry targets the same exhausted pool.

Systems insight: A robust scheduler diversifies compatible restart destinations, monitors interruption risk, and plans an on-demand fallback when spot capacity cannot meet the job's deadline or cost target. That fallback still depends on available capacity and the time needed to restore state.

8.8.4 Capacity reservation strategies

Committing to long-term reserved capacity for 100 percent of peak demand leaves expensive GPUs idle whenever traffic dips, while relying solely on spot or on-demand pools risks capacity shortages during product launches and sudden reclamations during deadline crunches. To avoid paying for idle reservations or starving urgent workloads, organizations that rely on cloud infrastructure (or that price on-premises capacity in economic terms) combine three tiers of compute capacity. Each tier trades cost and commitment against availability and interruption risk (table 8.4).

Table 8.4: Cloud Capacity Reservation Tiers: Capacity tiers trade cost and commitment against availability guarantees and interruption risk, giving the scheduler a price ladder for routing workloads by urgency and restart tolerance.

Capacity tier	Cost and commitment	Availability and interruption risk	Best-fit workloads
Reserved capacity	Capacity reservation; price discounts require a separate purchasing contract.	Capacity assurance within the reservation terms; startup still takes time.	Predictable continuous workloads: 24/7 inference, scheduled training, CI and regression testing, jobs under service-level agreements (SLAs).
On-demand capacity	Full price with no long-term commitment.	Admission depends on available capacity; provisioning and startup add delay.	Urgent retraining, interactive debugging, and short-lived jobs that do not justify spot management.
Spot capacity	Discounted interruptible capacity.	Availability varies with cloud demand, and workers can be reclaimed.	Restart-tolerant jobs: hyperparameter sweeps, ablations, checkpointed pretraining, and exploratory research.

The constructed portfolio in figure 8.14 divides capacity into a reserved tier covering the steady 60 percent baseline, an on-demand band between 60 and 80 percent, and an interruptible spot tier serving peaks up to 92 percent. An organization with 60 percent steady-state utilization and 40 percent burst demand might reserve capacity for 60 percent of peak need, use on-demand for bursts whose deadlines tolerate admission and startup delay, and use spot for restart-tolerant work. These shares illustrate a policy rather than establish an optimal portfolio. The scheduler must match each workload's deadline and recovery contract to the capacity tier. Reserved capacity can support the steady serving load or a training job's base allocation, while spot can support elastic replicas or independent experiments that tolerate reclamation. A peak in urgent demand cannot automatically be assigned to spot merely because it lies above the reserved band.

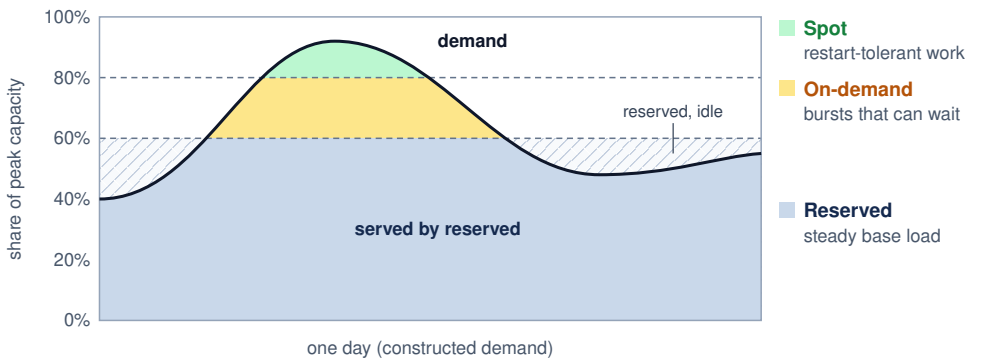

Figure 8.14: Capacity Reservation Tiers and Portfolio Optimization: This constructed portfolio combines reserved physical capacity, on-demand acquisitions, and interruptible workers. Reserved capacity covers the steady 60 percent base, on-demand the next 20 percent, and spot the peak; these shares and the demand curve are scenario assumptions. A discount commitment alone does not reserve capacity, and on-demand admission is not instantaneous. Schedulers map steady-state baseline demand to reserved capacity, diurnal spikes to on-demand instances, and restart-tolerant training to discounted spot pools. Hatching marks reserved capacity that is held while demand sits below it.

Systems Perspective 8.4: Spot vs. on-demand training

Consider training a 70B parameter model, a smaller workload than the chapter's 175B reference, on 512 GPUs for 14 days. The scenario fixes the event budget at 196 interruptions over the run and treats each as graceful, so no progress is replayed.

On-Demand: 512 GPUs $\times$ 336 hours $\times$ \$2/GPU-hour = \$344,064

Spot (65 percent discount, 5 percent checkpoint overhead, 14 interruptions per day of useful training with 30 minutes restart):

- Base cost: 512 GPUs $\times$ 336 hours $\times$ \$0.70/GPU-hour = \$120,422
- Checkpoint overhead: $5\% \times$ 336 hours = 16.8 hours
- Restart overhead: 196 interruptions $\times$ 0.5 hours = 98 hours
- Total time: 336 hours + 16.8 hours + 98 hours = 450.8 hours
- Total cost: 512 GPUs $\times$ 450.8 hours $\times$ \$0.70/GPU-hour = \$161,567

Savings: \$182,497 (53.0 percent), at the cost of approximately 34.2 percent longer wall-clock time.

These savings assume fault tolerance infrastructure is already in place. Without checkpointing, a single interruption forces restart from scratch, potentially wasting days of compute and negating all savings.

8.8.5 Cost-aware scheduling policies

Once a fleet spans reserved, on-demand, and spot tiers, a newly submitted job can often launch immediately on reserved A100 GPUs or wait two hours for spot H100 instances that, under illustrative assumptions, deliver twice the throughput at half the execution cost. Deciding whether those savings repay the two-hour delay and interruption risk turns the purchasing tiers into an online scheduling decision. Alongside completion time, utilization, and fairness, a cost-aware orchestrator evaluates three recurring trade-offs: waiting versus running immediately, preempting lower-priority work on reserved capacity (such as a hyperparameter sweep) versus routing a high-priority training run to more expensive on-demand instances, and replacing reclaimed spot instances with on-demand nodes versus shrinking elastically. Each choice requires the costs of waiting, recovery, and displaced work as well as the price of the new allocation.

These decisions require the scheduler to model both the time-value of computation (the organizational cost of a 2-hour delay in researcher productivity, product launch schedules, or competitive position) and the financial cost of different resource allocations. Delaying a production model retrain by 2 hours might violate a service level agreement (SLA), while the same delay to an exploratory experiment can postpone a research decision or block dependent work. A workload class's delay tolerance is therefore a policy input grounded in its dependencies and deadlines, not an assumption that exploratory work has no time cost.

Production systems can define **cost policies** per workload class that encode these trade-offs as configuration rather than requiring per-decision human judgment. One policy places steady inference demand on reserved capacity and maintains ready replicas to meet its latency target; the reservation alone does not guarantee application availability or latency. Large training runs can combine a reserved minimum viable allocation with interruptible elastic expansion, requesting on-demand replacements when a deadline justifies the cost and capacity is available. Exploratory workloads with flexible deadlines can use spot capacity while accepting interruptions and variable completion times. These assignments express the organization's priorities rather than a universal mapping from workload type to purchasing tier.

The scheduler does not implement the organization's financial accounting system. It consumes the resulting price and policy signals: the marginal price of each capacity tier, the reservation a project may claim, the interruption risk it may accept, and the deadline it must meet. That accounting system reports or bills each team for its usage (showback or chargeback) and attributes shared costs; the scheduler uses its outputs only to rank feasible scheduling actions. Lifecycle total cost of ownership (TCO) reaches the scheduler the same way, as per-resource opportunity costs and constraints that record whether capacity is owned, reserved, on demand, or interruptible, rather than as a scheduling mechanism.

↪ Tag-based cost attribution, shared-cost allocation, and chargeback are developed in section 12.9.7.2.

↪ The lifecycle TCO framework, which combines training, inference, data, and iteration costs, is developed in section 12.9.6.

Checkpoint 8.2: Cost-aware scheduling trade-offs

Before moving to ML-specific schedulers, review how the cost-optimization mechanisms interact:

- ☐ **Spot break-even**: identify the additional inputs needed to compute a break-even checkpoint interval for the 512-GPU, 14-day training job of the spot comparison above (65 percent spot discount, 14 interruptions per day, 30-minute restart overhead), including checkpoint duration, the interruption-rate clock, and lost-progress assumptions.
- ☐ **Double return**: explain why fault tolerance infrastructure has a "double return" for reliability and cost savings, and identify the minimum fault tolerance investment needed to make spot instances viable.
- ☐ **Reservation sizing**: starting from the constructed portfolio that reserves 60 percent of peak, identify the demand history, reservation terms, unit prices, and unused-capacity costs needed to decide whether raising the reserved share to 70 percent reduces total spending.

Spot pricing and reservation strategy reduce the cost of capacity, but they do not determine which competing job should receive it. Runtime uncertainty, unequal completion delays, and the cost of interrupting an iteration remain even when prices are known. Resolving them requires signals from the workload itself. A scheduler that tracks learning progress can optimize time-to-accuracy, the wall-clock time until a run reaches a target validation metric, rather than execution speed alone.

8.9 Custom ML Schedulers

A short experiment waiting behind a long training run can suffer even when every running job has an efficient placement. Improving that outcome requires a policy for deciding which job receives the next available resources. Custom ML schedulers use different signals to make that decision: consumed service can prioritize jobs without runtime estimates, iteration boundaries can make preemption cheaper, normalized finish time can expose unfair delay, and measured goodput can guide joint resource and batch-size changes. Only the last of these requires a model of learning progress; the others improve completion time, utilization, or fairness without predicting convergence.

8.9.1 Duration-agnostic scheduling (Tiresias)

When a researcher submits an exploratory training run without knowing whether it will fail in ten minutes or run for a week, scheduling policies that depend on upfront duration estimates (section 8.4.3) either mispack the queue or penalize inaccurate guesses. **Tiresias** (Gu et al. 2019) eliminates the requirement for duration estimates entirely. It uses a two-dimensional attained service[11] scheduler that makes priority decisions based on what a job has already consumed, not what it claims it will consume. Jobs accumulate "service" based on GPU-time consumed, with priority decreasing as service increases. The two dimensions are time (how long the job has been running) and resources (how many GPUs the job uses), capturing both elapsed duration and resource intensity.

A discretized version groups jobs into service bins (illustratively, less than 1 GPU-hour, 1 to 10 GPU-hours, 10 to 100 GPU-hours, greater than 100 GPU-hours), giving lower-service bins higher priority. When job durations are heavy-tailed, a job that has consumed little service is likely to be short, so favoring low-service jobs approximates shortest-remaining-processing-time ordering without any runtime estimate. New jobs can receive service before long-running jobs, but their eventual duration remains unknown, and a high-priority gang still needs a feasible allocation. Tiresias also promotes jobs that have waited too long, trading some average-completion-time benefit for protection against starvation. The priority rule therefore replaces runtime prediction with an observable signal and an explicit waiting-time policy; it does not guarantee that every short job starts promptly.

The Tiresias evaluation reports lower average job completion time (JCT) on its testbed and trace-driven workloads (Gu et al. 2019). The result depends on job arrivals, sizes, and the comparison

[11] **Attained service scheduling**: A discipline that uses consumed service to set priority when remaining work is unknown. Unlike shortest remaining processing time (SRPT), it cannot distinguish a newly arrived long job from a short one before either receives service. Its benefit depends on the workload distribution; protection against starvation remains necessary when new jobs repeatedly displace older ones.

policy. Prioritizing low-service jobs can delay long-running jobs, so average completion time must be evaluated alongside starvation prevention and the delay experienced by each job class.

8.9.2 Iteration-aware scheduling (Gandiva)

Preempting a job in the middle of a training step either discards the partial step or requires saving gigabytes of temporary forward-pass activations alongside the model weights. **Gandiva** (Xiao et al. 2018) avoids that mid-step state transfer by exploiting repetitive training iterations to profile jobs and time-slice GPUs at iteration boundaries. To give a GPU to another job, Gandiva suspends the running job at the end of a minibatch, copies its live GPU memory to host memory, and resumes another job on the device; switching back reverses the copy. Several jobs can therefore take turns on one GPU, and the same suspend-and-resume path lets Gandiva migrate a job to better-placed devices or pack compatible jobs onto one GPU when their combined memory fits. The boundary is the cheap place to switch because temporary activations have been released, so far less state must move than in the middle of a step. Weights, optimizer state, and runtime allocations remain, so each switch still costs a transfer proportional to that persistent state, and a job whose persistent state nearly fills the device cannot share it. The scheduler must measure memory demand, interference, and transition cost before admitting additional jobs.

Gandiva also implements **grow-shrink elasticity**, applying the allocation mechanism of section 8.7 to jobs that explicitly permit resizing. It expands eligible jobs into idle capacity and shrinks them when demand returns, using profiles to estimate which jobs benefit. The published implementation limits growth to one server and waits through an idle period before expanding. Its contribution is the connection between measured job behavior and resource control, rather than a claim that arbitrary training jobs can resize cheaply.

8.9.3 Finish-time fairness (Themis)

Equal GPU-time shares can produce unequal completion delays when workloads respond differently to placement. **Themis** (Mahajan et al. 2020) measures fairness in normalized finish time. Each application, which can contain related training jobs, compares its completion time in the shared cluster with its completion time on an isolated $1/N$ share of the cluster, where N is the number of applications contending for it. The ratio has the same form as the slowdown S_j defined with the scheduling objectives, except in its denominator. Themis divides by the completion time on a fair share of the cluster rather than by the isolated runtime on the job's own allocation, so a ratio of one means the application finished no later than its fair share would have allowed.

Themis targets the worst finish-time fairness ratio across applications. Its auction rounds offer available resources to the applications farthest from fair completion, while allowing placement efficiency to influence which specific GPUs each one receives within a round. With two applications whose current estimates are 1.2 and 2.0, the second is projected to finish twice as late as its fair share would allow, so it enters the round ahead of the first. The first application's ratio rises while it waits, and once it overtakes the second it moves to the front of later rounds. The signal is relative completion delay, not proximity to completion or GPU-hours already consumed.

Within a round, contending applications bid on candidate GPU subsets by reporting the finish-time fairness ratio they would reach on each subset. Because a communication-heavy job finishes much faster on a compact subset than on a scattered one while a communication-light job sees little difference, these bids expose placement sensitivity rather than monetary willingness to pay. The arbiter can then assign compact subsets to the lagging applications whose finish-time ratios gain the most from locality. This auction trades short-term allocation efficiency against the long-term fairness objective, making the quality of completion-time and placement estimates part of the scheduler's operating assumptions.

8.9.4 Adaptive resource allocation (Pollux)

Adding GPUs to a running job by increasing its global batch size raises throughput in examples per second. Past the critical batch size, however, each example contributes less optimization progress, so the extra GPUs can slow convergence rather than speed it up. **Pollux** (Qiao et al. 2021) addresses this coupling by extending resource adaptation to training hyperparameters. Where Gandiva changes only device allocations, Pollux jointly considers allocation, batch size, and learning-rate scaling

because a larger batch can use more devices efficiently while changing the amount of work needed to converge. The preferred configuration can therefore change with both cluster demand and the stage of training.

Pollux uses statistical goodput to connect system throughput, measured in training examples per second, with statistical efficiency, which measures progress per example relative to a reference batch size. The word goodput already carries two other meanings. The run-level goodput of section 1.5.4 discounts work lost to failures and restarts, and Google's ML productivity goodput multiplies program, runtime, and scheduling factors. Statistical goodput instead discounts each example by the optimization progress it contributes at the current batch size, and section 9.2.4 sets the three side by side. Statistical efficiency falls as the batch grows past the critical batch size of section 6.4.6, which is why more GPUs at a larger batch can raise throughput while lowering goodput. Pollux estimates how allocation and batch changes affect both quantities, then uses those estimates to guide resource sharing across jobs. Additional GPUs are valuable only when their throughput gain survives the corresponding change in statistical efficiency and the cost of changing the allocation.

The Pollux evaluation reports average job completion times 37 to 50 percent lower than the state-of-the-art deep learning schedulers it was compared against, which do not make the same combined adjustment of resources and batch size (Qiao et al. 2021). The number belongs to that testbed and trace. Its durable lesson is the coupling: an allocation can increase samples per second while reducing progress per sample. The scheduler must estimate both effects and validate completion time at the same target quality, rather than transferring a reported improvement to a different workload.

8.9.5 Scheduler comparison framework

Each research scheduler pursues a different objective, and each objective requires its own observations and controls (table 8.5). Consumed service can be obtained from resource accounting, while iteration timing and learning-progress estimates require closer runtime integration. The choice therefore turns on what the platform can observe and control, not on a ranking of scheduler sophistication.

Table 8.5: Custom ML Scheduler Comparison: A taxonomy of research schedulers based on the specific workload characteristic they exploit. Tiresias uses attained-service accounting, Gandiva profiles iterations, Themis estimates relative finish times and placement effects, and Pollux models goodput under allocation and batch-size changes.

Scheduler	Key insight	Scheduling signal	Optimization target	Integration requirement
Tiresias	Attained service informs priority	GPU-hours consumed	Minimize average JCT	Resource accounting and queue reprioritization
Gandiva	Iteration boundaries enable sharing	Iteration timing	Improve utilization	Runtime profiling and coordinated time slicing
Themis	Completion delay defines fairness	Normalized finish time	Improve worst fairness	Finish-time estimates and allocation auctions
Pollux	Batch size and allocation are coupled	Goodput	Maximize cluster goodput	Runtime resizing and batch-size adaptation

The common lesson is to match the policy signal to the scheduling objective. Attained service addresses runtime uncertainty and can work from accounting records. Iteration boundaries lower the cost of interrupting work but need reliable timing profiles. Normalized finish time measures unequal delay, and goodput couples resource allocation to learning progress, which requires the training loop to accept allocation and batch-size changes. Each added control creates a failure mode of its own: inaccurate profiles can increase interference, and an inaccurate learning-progress model can choose allocations that fail to improve convergence. A platform should therefore test the chosen policy against queue delay, starvation risk, and measured recovery cost before adopting it.

8.9.6 From research to production

A production cluster containing V100s alongside A100s and H100s, with both InfiniBand and Ethernet paths, gives the same GPU count very different performance across placements, while

live jobs crash, hang, or exceed declared memory needs. Because a research scheduler's evaluation establishes results on a specific testbed, workload trace, and set of runtime controls, a production fleet must test whether its own hardware and failure conditions preserve those gains. Deploying a policy therefore requires integrating its decisions with admission, failure recovery, and the existing orchestrator, then measuring its benefit under the fleet's actual workload.

Consider applying adaptive allocation to our 175B-parameter model training run. The runtime must first expose viable resize configurations and their costs. A checkpoint-and-restart path can require persisting and restoring large optimizer state, while other state-transfer mechanisms impose different costs. Repeated resizing can consume the expected throughput gain before the job benefits from its new allocation. The policy should therefore compare predicted additional progress over the expected allocation duration with the measured transition cost, retaining a stable placement when resizing does not repay that cost. The decision follows the workload and cluster demand, not a universal training phase.

The production decision is whether a specialized policy improves marginal allocation enough to justify its operational burden, judged by the measures this chapter has already built: queue delay and slowdown, useful GPU-hours, preemption overhead, and placement efficiency. Group admission settles whether a gang can start at all; the research schedulers answer the next question, which feasible group to admit, and each imposes different requirements on the runtime. Selecting a policy therefore means matching its objective and assumptions to the fleet rather than choosing a single scheduler as universally best.

Checkpoint 8.3: Custom scheduler design space

Consider the trade-offs between the four research schedulers:

- ☐ **Estimate-free scheduling**: identify the information Tiresias sacrifices by eliminating runtime estimates, and describe when that sacrifice hurts scheduling quality.
- ☐ **Time-slicing limits**: identify workloads for which Gandiva's iteration-boundary time-slicing fails, and explain why.
- ☐ **Job starvation**: explain how waiting changes an application's finish-time fairness ratio in Themis, and why this differs from favoring jobs that are nearly complete.
- ☐ **Goodput-model risk**: describe the failure mode when Pollux's goodput model is wrong.

The training policies just compared balance completion time, fairness, utilization, and learning progress. Serving adds a different deadline: capacity must be ready before an arriving request exhausts its latency budget. That requirement makes provisioning delay and burst headroom central to how the fleet allocates capacity to serving.

8.10 Serving Capacity in the Shared Fleet

When an e-commerce platform launches a major sale, requests can arrive faster than its serving replicas can process them. A training job may tolerate waiting in the admission queue; a serving request misses its deadline while additional capacity is still starting. To the fleet orchestrator, serving is therefore a claim on accelerators that must be honored quickly and that competes with training gangs for the same devices.

The orchestrator's unit is the serving replica, a ready execution group with a loaded model that may occupy one accelerator, a device partition, or a multi-accelerator gang. A 175B-parameter model may make an eight-GPU tensor-parallel group the schedulable replica. Archetype B (DLRM at Scale) imposes the same constraint through its embedding tables: a replica whose tables are sharded across several accelerators can serve only when every shard is placed. Each replica has a measured request-processing capacity at a specified workload mix and latency target, and that measured capacity is the quantity the orchestrator uses to size and place replicas. It decides how many such groups exist, where they reside, how they are protected from interference, and which lower-priority allocations yield when serving demand grows. Language models add a memory dimension to the claim. Each in-progress response retains per-token attention state, the key-value (KV) cache, whose

footprint grows with context length and concurrency even though the weights stay fixed, so a replica can exhaust its memory before it exhausts its compute.

How many replicas to run at each moment is a control problem with its own signals and delays. Accelerator utilization alone does not identify the capacity needed to meet a request-latency target, and loading large models can delay a new replica's readiness. The controller must combine demand signals with measured startup time and ready reserve capacity; section 10.13 develops that control loop. For the orchestrator, the immediate consequence is that reclaiming a device from training does not instantly make it available for serving. Three fleet-level questions follow, namely how serving workloads share devices without unacceptable interference, when sharing pays for itself, and where the boundary between training and serving pools should sit.

8.10.1 Resource isolation

When multiple inference workloads share physical hardware, noisy-neighbor contention[12] arises whenever one workload's resource consumption degrades another's execution speed. On GPUs, shared execution units can limit concurrent kernels, while saturated HBM interfaces constrain memory-bound work. Contending working sets can evict lines from the shared L2 cache, and simultaneous host-to-device transfers can compete for PCIe bandwidth. Aggregate power and temperature limits can also reduce operating clocks. These mechanisms connect the placement decision to the request's critical path. A placement that fits every model in memory can still miss its service-level objective if sharing extends execution or queueing enough to exhaust the latency budget.

Multi-Instance GPU (MIG)[13] closes several of these channels in hardware by partitioning A100 (80 GB) and H100 (80 GB) GPUs into hardware-isolated instances. Unlike software-based GPU sharing, such as time-slicing that context-switches several processes over one whole device, MIG dedicates memory, cache, and compute resources to each instance, reducing interference through those partitioned resources. The A100 MIG profiles in table 8.6 show the scheduling consequence. A single physical A100 becomes several fixed-size resources that the device plugin of section 8.5.2 can expose as independently schedulable accelerators.

Table 8.6: A100 MIG Profiles: Multi-Instance GPU configurations trade GPU memory and compute resources for increased sharing density. Each profile provides hardware-isolated resources, making MIG suitable for multi-tenant inference clusters where workloads from different teams or customers must not interfere. The streaming multiprocessor (SM) count determines compute throughput within each instance.

MIG profile	GPU memory	SM count	Typical workload
1g.10gb	10 GB	14 SMs	Small inference
2g.20gb	20 GB	28 SMs	Medium inference
3g.40gb	40 GB	42 SMs	Large inference
7g.80gb	80 GB	98 SMs	Training

MIG therefore sits at the strong resource-isolation end of this design space. The operational implication is that hardware partitioning reduces cross-workload interference but turns partition choice into a placement constraint. MIG is available on supported MIG-capable NVIDIA data-center GPUs, profiles may require draining affected workloads before reconfiguration, and fixed partition sizes may not match workload requirements. The memory-fit and sharing examples take each labeled device or profile capacity as a conservative admission budget of that many decimal gigabytes, so an instance admits a working set no larger than its label. A workload that needs 15 GB of GPU memory leaves 5 GB of a 20 GB MIG instance budget unused or exceeds a 10 GB instance budget.

Software sharing offers different controls over execution, memory, and fault isolation; these do not form a single performance-versus-security scale. At the simplest level, **CUDA time-slicing** allows multiple processes to share a GPU by context-switching the compute resources. While flexible, it offers weak performance isolation because contexts compete for the entire device over time. NVIDIA's Multi-Process Service (MPS)[14] allows kernels and memory copies from different processes to overlap, reducing context-switch overhead and improving utilization when one process

12 **Noisy neighbor**: In GPU clusters, interference arises when one job's bursty network traffic or memory bus utilization slows down a co-located job on the same node. In synchronous training, a delayed flow can extend collective completion when it lies on the communication critical path, propagating one tenant's interference across a worker group.

13 **Multi-Instance GPU (MIG)**: Introduced with the A100 (2020), MIG partitions a single GPU into up to 7 hardware-isolated instances with dedicated memory controllers, L2 cache slices, and compute units (NVIDIA Corporation 2020). The isolation is enforced at the hardware level, not by time-slicing, reducing noisy-neighbor interference relative to software-only sharing. Changing the partition layout requires stopping affected GPU workloads; the scheduler must include that disruption when adapting profiles to a changing workload mix.

14 **CUDA MPS (Multi-Process Service)**: The shared-resource mode used in this example does not reserve an independent hardware slice for each client, and a deployment must also account for the GPU fault domain its clients share (NVIDIA 2026d).

cannot fill the device (NVIDIA 2026d). In the shared-resource MPS configuration considered here, Volta and newer clients have distinct GPU virtual address spaces but compete for execution and memory-system resources. Additional partitioning controls depend on the hardware, software version, and selected mode, and those configured limits, not address-space separation, determine how much performance and fault isolation MPS provides.

GPU memory accounting must still prevent one model from consuming capacity reserved for another. A distinct virtual address space blocks direct addressing of another client's allocations, but it does not create a fixed memory-capacity partition. Without enforceable per-client limits, a memory leak or unexpectedly large batch can exhaust shared device memory and disrupt colocated workloads. MPS therefore fits mutually trusting workloads whose aggregate memory footprint and interference have been measured; MIG or full-device assignment fits stronger trust and predictability requirements. At the host level, CPU core pinning and NUMA-aware placement complement these device controls by keeping inference processes off contended cores and away from remote memory, which shortens tail latency whenever those host paths lie on the request's critical path.

The GPU mechanisms form a spectrum from flexible software sharing (time-slicing and MPS) through hardware partitioning (MIG) to exclusive assignment, and each point gives up packing efficiency for isolation (table 8.7). For serving, the choice turns on interference and memory accounting; when colocated workloads belong to different teams, it also becomes a trust-boundary decision, which section 8.11.3 takes up with the rest of multi-tenant security.

Table 8.7: GPU Virtualization Isolation Spectrum: GPU sharing mechanisms trade device-resource isolation against packing efficiency. External tenant isolation additionally depends on the host, driver, identity, network, and storage security boundaries.

Option	Isolation strength	Efficiency and flexibility	Suitable workload or trust boundary
Time-slicing	Separate contexts; shared performance and device fault domains.	Flexible; no spatial reservation.	Trusted, bursty jobs.
MPS	Separate address spaces; shared performance and GPU fault domains.	Concurrent kernels; configurable limits.	Mutually trusting jobs with measured demand.
MIG	Hardware-partitioned memory, cache, and compute.	Predictable; fixed profiles.	SLO-sensitive work; external tenants also need platform security controls.
Full device passthrough	Exclusive access to the GPU.	Strongest device isolation; lowest packing efficiency.	Jobs that saturate the device or forbid sharing.

8.10.2 GPU sharing economics

The decision to co-locate multiple models on a single GPU trades potential capacity savings against interference. Choosing **exclusive access** avoids competition from another device tenant but can leave capacity unused when one workload cannot fill the device. Production evidence suggests such stranded capacity is common. On NERSC's Perlmutter supercomputer, Li et al. (2023) found that 50 percent of jobs used less than 25 percent of their allocated GPU memory. Moving to **aggressive sharing** (via MPS or time-slicing) can admit more work, provided contention does not erase the gain or violate latency objectives. Exclusive device assignment still depends on the surrounding host, driver, identity, network, and storage controls for tenant security. The scheduler must therefore evaluate memory fit and measured service capacity together.

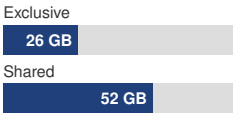

Sharing packs two tenants, cutting the unused memory budget from 67.5 percent to 35 percent.

Using the same conservative admission budget, consider an A100 GPU with a 80 GB memory budget serving a 7B parameter model. The model weights in FP16 consume approximately 14 GB. With the scenario's assumed KV cache and activation overhead, the total runtime footprint is roughly 26 GB. Exclusive access leaves 54 GB (67.5 percent) of the memory budget unused at this load. That budget is reclaimable only to the extent that it exceeds required runtime and burst headroom. A MIG partition (for example, a `3g.40gb` profile) leaves roughly 35 percent of the assigned partition budget unused; the remaining device capacity is subject to valid partition combinations. Packing two such models onto the same 80 GB device through MPS consumes 52 GB, leaving 35 percent of the device memory budget unused under the same footprint assumption. Reducing the device

count also requires enough compute and bandwidth headroom to preserve throughput and latency under concurrent demand; fitting the models in memory is only the first admission condition.

This density introduces interference. When two workloads share execution and memory resources through MPS, their demands can overlap on SMs and memory bandwidth. A compute-bound model paired with a memory-bound model is a candidate for complementary placement because, in isolation, each saturates a different resource, but the bottleneck labels describe isolated behavior, not a guarantee of compatibility. Both workloads still use multiple resources, and concurrent demand can change the limiting resource. The scheduler must measure the pair's throughput and tail latency across its expected load range before treating memory density as usable capacity.

For a heterogeneous fleet, this logic dictates a bifurcated strategy. Our 175B parameter model requires about 350 GB for FP16 weights alone; with the 80 GB per-device admission budget that means at least 5 GPUs before KV cache, activation buffers, runtime workspace, and tensor-parallel granularity are included. In practice, an 8-GPU tensor-parallel replica is the cleaner scheduling unit and is a candidate for exclusive access. Conversely, the dozens of 1B to 7B parameter support models, such as toxicity classifiers, query embedding encoders, and auxiliary generation models, are candidates for sharing. Packing these smaller models onto shared utility nodes frees the high-performance nodes for foundation-model replicas.

Napkin Math 8.4: GPU sharing ROI

Consider a fleet serving 100 distinct 7B models (26 GB footprint each) and 10 distinct 175B models (350 GB footprint each). Assume load testing has established that each proposed co-resident pair meets its throughput and tail-latency objectives with the required headroom and isolation. The memory-utilization percentages use the same conservative admission budgets. The example computes the resulting capacity reduction under the load-testing assumption; memory fit alone would not establish it.

Strategy A (exclusive access, one model per GPU): The 7B models require 100 A100 GPUs (80 GB each), achieving only 32.5 percent memory utilization (100 $\times$ 26 GB used out of 8000 GB total). The 175B models require 80 GPUs (10 models $\times$ 8 GPUs each for tensor parallelism). Total fleet: 180 GPUs.

Strategy B (MPS for small models, exclusive for large): Pack 2 models per A100 for the 7B models (2 $\times$ 26 GB = 52 GB, well within 80 GB), requiring only 50 GPUs at 65 percent memory utilization. The 175B models remain unchanged at 80 GPUs. Total fleet: 130 GPUs.

Under the load-testing assumption, sharing reduces the required fleet size by 27.8 percent (50 fewer GPUs). At $2/GPU-hour, the released capacity is worth approximately $876,000 annually at the scenario rate. It becomes a cash saving when the organization releases billed resources or avoids purchasing that capacity.

8.10.3 The training-serving resource boundary

Separate training and serving pools simplify capacity ownership but can leave useful work queued beside idle devices. Consider a serving pool reserved for peak traffic while a training pool remains backlogged. This **static partitioning** protects serving capacity from training demand, but unused serving resources cannot help the training queue. **Dynamic sharing** allows that borrowing only when the orchestrator can return ready serving capacity before it is needed. The mode-switching cost determines how much capacity can safely move between the pools.

Repurposing a GPU from training to serving can require saving training state, releasing the old allocation, launching the inference runtime, loading weights, and warming the replica. The orchestrator must measure the complete path to readiness, including any overlap between these operations. If that path takes minutes, reclaiming training capacity cannot cover a burst that exhausts the serving reserve in seconds. Forecasts can schedule transfers ahead of predictable demand, while already-ready headroom covers shorter bursts; the required reserve depends on forecast error and measured reclaim time.

Consider a daily traffic pattern with low demand at 3 AM, growth at 8 AM, a peak at 2 PM, and a decline at 10 PM. A forecast with a 30-minute lead allows the orchestrator to begin reclaiming

training allocations at 7:30 AM, provided measured replica startup fits within that lead time. As traffic falls, it drains surplus replicas before returning their devices to training. Pricing that policy against a static split needs three quantities: what serving reserves, what it uses on average, and how much of the difference training can absorb without eroding the surge buffer. The example below assumes its surge buffer suffices; a production policy must size the buffer against its own burst distribution and forecast errors.

Napkin Math 8.5: The diurnal GPU shift

Consider a fleet of 1,024 GPUs managed under two regimes:

Scenario A (static partitioning): The serving pool reserves 400 GPUs. That reservation is the forecast peak of 283 GPUs plus a 117-GPU surge buffer that absorbs forecast error and bursts shorter than a replica's startup time. Average usage is only 150 GPUs (37.5 percent utilization). The training pool takes the rest of the fleet after the serving reservation and 6 GPUs held back as maintenance headroom, which leaves 618 GPUs, and its backlog keeps them at 100 percent utilization. Result: 150 GPUs + 618 GPUs = 768 active GPUs. Fleet utilization: 75 percent. The maintenance headroom stays unallocated in both scenarios.

Scenario B (dynamic sharing): Serving holds forecast demand (average 150 GPUs) plus the same 117-GPU surge buffer. Training reclaims the rest of the serving reservation, which is the unused 250 GPUs on average minus the buffer and averages 133 GPUs over the course of the day (the medium-blue band in figure 8.15). Because each role switch costs minutes, the orchestrator moves this capacity on the forecast rather than chasing individual traffic spikes. Result: serving (150 GPUs) + training (618 GPUs + 133 GPUs) = 901 active GPUs (the second training term is reclaimed serving capacity). Fleet utilization: 88 percent.

Dynamic sharing recovers 133 GPUs of effective capacity. At $2/GPU-hour, this creates approximately $2.3M per year in value without purchasing a single additional accelerator.

Safe sharing requires both a measured interference budget and an enforceable right to reclaim capacity. Isolation mechanisms address the first requirement; they do not decide which team owns a device or when a borrower must release it. That second requirement is an allocation contract across teams, and it applies to every pool the orchestrator manages, serving and training alike.

8.11 Multi-Tenancy and Quotas

When the computer vision team hoards 500 GPUs for a week while the natural language processing team sits blocked trying to push a critical bug fix to production, shared clusters devolve into a tragedy of the commons. Teams with the most aggressive job submission rates, the largest jobs, or the most persistent resubmission scripts consume disproportionate resources, while teams with more modest or intermittent needs find the cluster perpetually occupied. The result is organizational friction and underinvestment in slower-moving but potentially higher-value projects whose teams cannot compete for capacity on submission volume alone. A static serving reservation faces the same pressure from the other direction, since a training team that sees idle serving devices will want them.

Multi-tenancy and quotas turn organizational priorities into rules for resource allocation, borrowing, and reclamation that the scheduler can enforce. A quota defines an entitlement or a limit. Whether that entitlement becomes prompt, productive access depends on borrowing that lends idle guarantees to other teams and on reclamation that returns them within a bounded time, both developed below. Accounting that separates the capacity a team holds from the capacity it uses well belongs with utilization debugging in section 8.12.1.

8.11.1 Hierarchical fair-share

Assigning fixed GPU quotas at the namespace or project level (for example, 500 GPUs to Team A, 300 to Team B, and 200 to Team C) gives each team a clear entitlement, but leads to systematic underutilization whenever workloads fluctuate. If Team A's 500-GPU allocation sits 50 percent idle during a quiet period while Team B's queue overflows with urgent work, rigid quotas leave 250

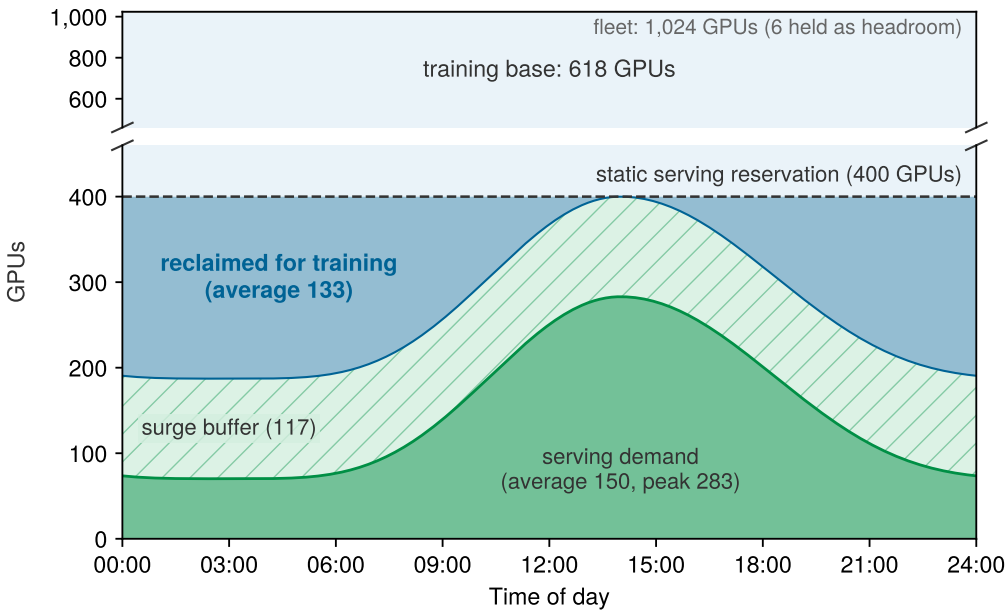

Figure 8.15: Training vs. Serving Diurnal GPU Capacity Shift: One day on the 1,024 GPUs fleet of the diurnal GPU shift notebook, drawn from the same numbers. Serving demand (green) averages 150 GPUs and peaks at 283 GPUs; the orchestrator keeps a 117-GPU surge buffer above the forecast, so the serving side never exceeds the 400 GPUs a static partition would reserve (dashed line). Everything between the buffer and that line returns to training, averaging 133 GPUs over the day, on top of the 618 GPUs training base. The remaining 6 GPUs of maintenance headroom at the top of the fleet are too thin to draw and are named in the upper panel. The shape of the daily curve is illustrative; its average and peak are the notebook's.

GPUs wasted that Team B is forbidden to use. Allocating smaller static quotas to each team creates the opposite failure: teams are blocked from running legitimate work during peak periods even when the cluster has spare capacity overall.

Hierarchical quotas, which Slurm expresses through its account hierarchy and fair-share accounting (equation 8.1) and Kueue through ClusterQueues grouped into borrowing cohorts, separate a team's guaranteed share from its maximum permitted allocation. In equation 8.4, Q_{team} denotes a team's maximum burst cap rather than its guarantee, $Q_{\text{department}}$ is the department's total capacity limit, and $U_{\text{allocated}}$ is each sibling team's current active allocation. Assuming those allocations remain within the department limit, the smaller of the team's burst cap and the department's currently unallocated capacity sets the team's admission ceiling before any reclamation:

$$Q_{\text{effective}} = \min\left(Q_{\text{team}},\ Q_{\text{department}} - \sum_{\text{other teams}} U_{\text{allocated}}\right) \tag{8.4}$$

The second term inside the minimum is what lends idle guarantees across sibling teams. If Team A uses only 60 percent of its 500-GPU guarantee, its active allocation $U_{\text{allocated}}$ drops below its guarantee and leaves the remaining 200 GPUs in $Q_{\text{department}} - \sum_{\text{other teams}} U_{\text{allocated}}$ for sibling teams to claim, up to each borrower's own cap Q_{team}. This ceiling bounds what a team may hold at admission, and it encodes neither fair-share weights nor guarantees. When several teams contend for that spare capacity at once, the scheduler still applies its fairness policy, the dominant-resource calculation of section 8.4.1, to the feasible gangs that fit under each ceiling.

8.11.2 Burst capacity and quota borrowing

Lending turns unused guarantees into burst capacity. Suppose a small team holds a guarantee of 64 GPUs but needs 256 for a weekend-long training run while a sibling team's guaranteed share sits idle. The small team may borrow idle whole devices up to its maximum quota. Those devices remain part of physical cluster capacity rather than becoming virtual capacity, and they carry lower

preemption priority than owned resources. When the sibling team submits work that needs its guarantee, the scheduler reclaims the borrowed devices by preempting or shrinking the borrower's opportunistic jobs.

Borrowing resolves that static-quota dilemma only when the lent capacity remains reclaimable within the platform's response-time objective, safety headroom, and current physical capacity. A team entitled to 500 GPUs can reclaim borrowed devices according to that contract, but the scheduler must not promise an immediate 500-GPU allocation when failures, nonpreemptible reservations, or recovery headroom reduce reclaimable capacity.

Burst capacity therefore raises the maximum a team may hold, while its guarantee stays fixed. The small team above may hold up to 256 GPUs, but the platform promises only its 64-GPU guarantee concurrently with every other team's. The sum of guarantees that the platform promises to honor concurrently must fit reclaimable physical capacity after nonpreemptible reservations and safety headroom. Teams may advertise higher maximums or fair-share weights, the target shares that set priority in equation 8.1, because borrowed jobs yield when guaranteed demand returns, but those values are not guarantees.

Reclamation couples quota policy to the checkpoint infrastructure. The admission controller separates guaranteed jobs from opportunistic jobs. Guaranteed jobs run on the team's owned quota and face preemption only from higher-priority work. Opportunistic jobs run on borrowed capacity, yield first, and must checkpoint often enough that reclamation wastes little work, since every forced reclaim pays the preemption overhead of section 8.8.1. Elastic jobs soften reclamation further through the graceful preemption of section 8.7.2, shrinking one replica at a time toward their minimum viable allocation rather than terminating.

One policy for quota borrowing uses four tiers. Production serving (Priority 0) holds capacity protected from routine preemption. Production training (Priority 1) yields only under a defined serving-failover policy. Interactive and debug sessions (Priority 2) receive short execution limits to preserve iteration speed without indefinite resource occupation. Batch and research tasks (Priority 3) use surplus capacity and yield under their reclamation contract. These numbers label this example's tiers; they are not Kubernetes priority values, where larger integers indicate higher priority. The hierarchy establishes precedence, but deadline protection still requires feasible capacity and bounded reclamation time, while lower-priority work needs an explicit starvation policy.

As the over-subscription discussion in section 8.4.2 showed, a partly idle device is still a whole allocated device, so the safe burst limit follows from reclaimability rather than average SM utilization. Fractional sharing is a separate mechanism that requires an explicit MIG, MPS, or time-slicing contract. The scheduler must measure how quickly borrowed jobs can checkpoint, shrink, or terminate, and it must reserve the I/O and network bandwidth required for that transition. Only then can it bound the amount of opportunistic capacity that is safe to lend while still honoring the guarantee's response-time objective. Until reclamation completes, a borrowed job that still runs beside the owner's returning work is a noisy neighbor on the shared network and storage paths (section 8.10.1). Isolation and the reclaim deadline must together bound how much the borrower can slow the owner.

8.11.3 Security and namespace isolation

Teams that share a fleet must also be kept from observing or altering one another's model weights, training data, and inference traffic. A namespace provides only an administrative boundary that host, device, network, and storage controls must enforce. Network policy must cover the collective paths an admitted gang actually uses, and host networking moves the enforcement point off the pod network, so each path needs its own verification.

↩ Fabric-level isolation, through virtual lanes and traffic classes that keep one job's bursts from stretching another job's collectives, is established in section 3.8.2.

Across teams, the isolation spectrum of table 8.7 becomes a trust-boundary choice, because colocated workloads now belong to owners with different exposure to one another. Side-channel attacks on shared GPUs are a documented vulnerability; by monitoring contention on shared caches or memory controllers, a malicious tenant can infer the architecture or even data properties of a co-resident model. In sensitive environments, MIG or dedicated devices can strengthen device-resource isolation, but neither creates a complete external-tenant security boundary by itself. A multi-tenant platform combines device partitioning with VM or container isolation, configuration of the input-output memory management unit (IOMMU) that limits which host memory a device can

reach, driver and runtime hardening, authentication and authorization, and network and storage controls. Single-organization workloads with a narrower trust boundary can use MPS or time-slicing when weaker fault and performance isolation is acceptable.

↪ The attack surface of shared accelerators, and which sharing mechanisms leave co-resident tenants observable to one another, is developed in section 13.6.3.

8.11.4 Quota governance and organizational dynamics

Even when a scheduler enforces quotas accurately, a guarantee set six months ago can leave 60 percent of a team's reserved GPUs underused after its flagship model ships while another team's urgent launch queue backs up. While the scheduler manages immediate resource assignments, **quota governance** revises the entitlements and limits those assignments must honor. Periodic **quota review cycles** compare observed use, forecast demand, and the cost of meeting each team's capacity promise. A team using only 40 percent of a 128-GPU guarantee does not necessarily block other teams if unused devices can be borrowed. The governance question is whether the guarantee remains justified and whether its reclamation contract lets the fleet use spare capacity productively between demand peaks.

The central tension in quota management is the **hoarding problem**, in which teams retain idle reservations to avoid queuing again when work becomes ready. A workflow that pauses policy training while awaiting human annotations or a training run delayed by unfinished tokenization can leave its reserved accelerators unused, reflecting dependencies in particular workflows rather than an unavoidable idle phase of all training. Utilization-based reclamation can release idle physical allocations to borrowers, provided the platform preserves any promised return time; revising the team's longer-term guarantee instead requires comparing expected demand with observed use. Clear return contracts reduce the incentive to hold idle devices without making low activity alone grounds for breaking a capacity promise.

Quota accounting should remain close to the scheduling decision. Every admitted job needs an owner, project, workload class, and preemption contract so the scheduler can apply quota and recovery policy consistently, and Kubernetes labels and Slurm accounts carry this scheduling identity. GPU-hours charged against a reservation, borrowed share, or preemption class then make the opportunity cost of hoarding visible and let the control plane reclaim idle guarantees.

> Example 8.1: Quota hoarding
>
> **Scenario**: A research team reserves a static block of 500 GPUs for an upcoming model refresh.
>
> **Diagnosis**: Upstream data delays postpone training, but the team retains the reservation to ensure future availability. The GPUs sit idle for 3 weeks, leaving $504,000 of capacity value unused while other teams' queues back up.
>
> **Systems lesson**: An idle reservation should trigger a review of lending eligibility and future demand. The orchestrator can admit opportunistic borrowers under a bounded reclamation contract without changing the owner's guarantee. Reducing that guarantee is a separate governance decision; lending owned capacity does not turn it into cloud spot capacity.

A "use it or lose it" reclamation policy addresses the most visible symptom of quota hoarding. Lending idle capacity to another team, however, requires the system to distinguish borrowed capacity that is preemptible on short notice from guaranteed reservations protected by SLOs. Defining borrowing, preemption, and guaranteed tiers as first-class policy primitives is the core multi-tenancy design decision.

> Checkpoint 8.4: Multi-tenancy design decisions
>
> Consider a 2,000-GPU cluster shared between a research team (60 percent allocation) and a production team (40 percent allocation):
>
> - ☐ **Idle reclamation**: decide whether the production team can borrow the research team's allocation while half of it sits idle, and describe the preemption behavior when the research team submits a large job.

- ☐ **Priority classes**: configure priority classes for a production inference workload that needs 100 GPUs with guaranteed latency SLOs and a research training job that can tolerate preemption.
- ☐ **Sharing contract**: explain why 65 percent average SM utilization does not justify whole-device overcommitment, then identify which measured resource and isolation constraints would make statistical multiplexing safe.

Even with isolation, quotas, and dynamic sharing in place, a cluster can report low utilization beside a full queue. Top-level dashboards rarely explain why pending jobs cannot claim the free devices. Locating that loss requires systematic utilization debugging across hardware pools, scheduler constraints, and user request habits.

8.12 Debugging Cluster Utilization

A cluster dashboard shows average GPU utilization stuck at 60 percent, yet developers complain about three-day queue times for small debugging jobs. Hundreds of GPUs sit idle, but no jobs are starting. Explaining the gap requires tracing the interactions among hardware topology, scheduler policies, and user behavior.

Part of the paradox is the queueing arithmetic of notebook 8.1. The $\rho_{\text{serv}}/(1-\rho_{\text{serv}})$ factor grows without bound as offered load approaches service capacity, so mean wait rises gently through moderate load and then steeply above 80 percent for small, mixed, and large workloads alike (figure 8.16, plotted on a logarithmic scale). The three curves differ only in mean job duration, as their labels show, and all use the notebook's exponential reference case ($c_s = 1$), so they are the gentlest version of the effect. A mixed ML queue with the notebook's representative $c_s = 3$ multiplies every curve by the larger variability factor, shifting each one upward without changing where it bends. The curves also describe offered service load rather than allocated or kernel-active GPU percentages, and their 75 percent headroom marker is an illustrative choice rather than a universal cluster threshold. Queueing explains why a busy fleet waits. A fleet whose queue is full while devices sit idle is waiting for a different reason, and finding it requires looking past the dashboard.

8.12.1 Allocated, active, and productive utilization

The 64 GPUs that sat idle for a weekend because a researcher forgot to cancel a reservation represent invisible waste in a multi-tenant cloud fleet. In a dedicated cluster, such waste is at least visible to the team that owns the hardware. In a shared environment, the cost is hidden and compounds across thousands of jobs. Effective governance requires moving beyond a single utilization number to a tiered accounting model that distinguishes between what was allocated, what was used, and what was useful.

Three tiers of measurement expose this gap between allocated hardware and actual progress. Allocated capacity counts resources held by complete allocations, so devices assigned to a gang that still waits for its remaining members do not count. A partial gang can run no step, so those held devices are a scheduling loss below the allocated tier rather than allocated capacity. Active utilization measures hardware activity under a stated counter definition, which can include communication kernels or work later replayed. Productive utilization requires a workload-level measure of retained training progress or serving work that meets its objective, normalized to the job's own fault-free step rate. For example, a team allocated 500 GPUs might average 350 active-device equivalents (70 percent activity) while delivering 280 device-equivalents of useful work at its fault-free rate (56 percent of its allocation). These are time-averaged equivalents, not necessarily distinct sets of active and productive devices. The gap motivates profiling; it does not identify its cause or imply that necessary communication and recovery can all be eliminated. The allocated and productive tiers also serve as the boundaries of two goodput factors that section 9.2.3 measures. Scheduling goodput is the share of offered capacity held by complete allocations, and runtime goodput is the share of allocated time that becomes retained progress.

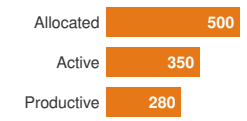

Allocated, active, and productive capacity in GPU equivalents, each tier counting less than the one above.

This data feeds the utilization dashboard, the cluster operator's primary instrument for fleet health. This view must synthesize per-team utilization against quotas, queue depth by priority class, and the **fragmentation index**, a measure of free GPUs that cannot be allocated due to topology

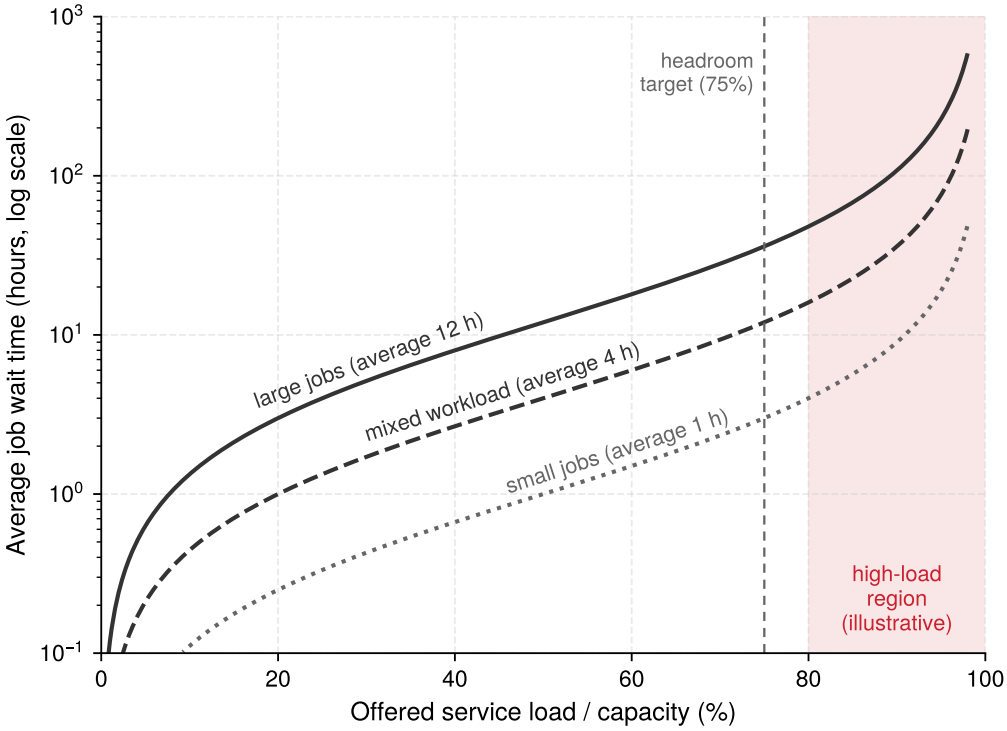

Figure 8.16: The Utilization Paradox: As cluster utilization approaches saturation (100 percent), job wait times increase hyperbolically and diverge near saturation, not linearly. The shaded region above 80 percent and the dashed 75 percent target illustrate headroom choices in this toy queue model, not measured GPU-utilization thresholds. Real queue delay also depends on job sizes, gang constraints, and arrival variability.

constraints. It must also track preemption rates, spot interruptions, checkpoint traffic, and the recovery overhead they induce. These measurements let the scheduler distinguish idle capacity it can reclaim from allocated capacity whose apparent idleness reflects communication, I/O, or a rigid gang dependency.

Automated anomaly detection turns these metrics into investigation triggers. A sudden spike in a team's consumption to 3× its historical norm might reflect a runaway submission script or an intentional experiment. A growing queue can indicate insufficient capacity, but it can also reflect quota limits, fragmented placements, or a stalled dependency. A sudden power drop can accompany a shared storage stall, though job completions and changes in workload mix can produce a similar signal. Operators must correlate these alerts with admission decisions and job progress before changing capacity or policy.

Systems Perspective 8.5: The three utilization metrics

Each tier narrows the investigation. Low allocated utilization calls for checking demand, quotas, device health, and placement feasibility. Low active utilization on allocated devices calls for tracing input supply, communication waits, and runtime execution. Low productive utilization despite busy kernels calls for comparing retained progress with replay, synchronization, and other overhead. The metric selects the next evidence to collect; it does not by itself select the remedy.

8.12.2 Policy-infrastructure mismatch

When a cluster adds a new generation of high-memory GPUs alongside older pools, job templates written for the newest hardware can leave the older pools sitting idle while dozens of jobs queue behind the newest partition. Diagnosing this **policy-infrastructure mismatch** requires comparing scheduler state with physical capability and inspecting job requests against each hardware pool. A full queue and idle accelerators can coexist when the scheduler is honoring constraints that the top-level dashboard has aggregated away. The cluster at the center of example 8.2 runs the heterogeneous hardware in table 8.8, accumulated across three procurement cycles.

Table 8.8: Cluster Hardware Inventory: Three procurement generations create distinct resource pools with different capabilities. In this inventory, A100 nodes have NVLink and the V100 pool uses PCIe. That difference changes tensor-parallel communication cost; it does not make tensor parallelism impossible on V100 devices.

GPU type	Count	Memory	Interconnect	Nodes
A100-80 GB	400 GPUs	80 GB HBM2e	NVLink (600 GB/s)	50 nodes × 8 GPUs
A100-40 GB	424 GPUs	40 GB HBM2e	NVLink (600 GB/s)	53 nodes × 8 GPUs
V100-32 GB	200 GPUs	32 GB HBM2	PCIe Gen3 (15.8 GB/s)	50 nodes × 4 GPUs

Example 8.2: Debugging low GPU utilization

Scenario: A 1,024-GPU cluster reports 60 percent average utilization despite over 50 pending jobs in queue.

Diagnosis: Over-specification (users copying A100-80 GB templates for 35 GB memory jobs), pool fragmentation, and legacy V100 avoidance leave A100-40 GB and V100 nodes idle while large jobs wait for A100-80 GB allocations. The scheduler correctly enforced policies written for a homogeneous fleet that no longer matched the heterogeneous infrastructure.

Systems lesson: Surface-level diagnosis often mistakes policy-infrastructure mismatches for scheduling algorithm bugs. Updating scheduler policies to support automatic capability matching, dynamic memory profiling, and backfilling raises cluster utilization to 84 percent, recovering $7M in annual compute capacity. Table 8.9 breaks that gain out pool by pool, using the allocated and active tiers of section 8.12.1.

Table 8.9: Utilization Recovery Results: Before the fix, A100-80 GB GPUs were almost fully allocated but less than fully active, while strict job templates stranded the other pools. Pool rows report allocated utilization, and for the A100-80 GB pool also the active share of its allocated GPUs; the A100-80 GB change is in that active share, the others in allocation. The cluster-wide row is the fleet's separately observed average utilization, so the pool rows are diagnostic signals that locate the loss rather than terms that add up to the total.

Pool	Before	After	Change
A100-80 GB	95% allocated, 72% active	89% allocated, 94% active	+22 pp
A100-40 GB	64% allocated	88% allocated	+24 pp
V100	0% allocated	71% allocated	+71 pp
Cluster-wide	**60%**	**84%**	**+24 pp**

For this capacity valuation, assume that each active GPU-hour delivers the same retained useful work before and after the policy change, including across hardware pools. Under that simplifying assumption, the observed utilization gain is equivalent to approximately 246 fully active GPUs

of productive work (table 8.9). A heterogeneous production fleet must validate that conversion against workload progress and measured throughput in each pool. As at the start of the chapter, procurement planning compares the gain with the hardware it replaces, here 410 additional raw GPUs, because new devices under the old policies would also run at only 60 percent utilization. At $2/GPU-hour, that avoided purchase represents approximately $7M in annual recovered value, achieved through policy changes requiring zero additional hardware investment.

8.12.3 Common utilization anti-patterns

Fixing policy-infrastructure mismatches gets queued jobs onto GPUs, but a dashboard reporting 85 percent allocation still says only how much capacity is assigned, not whether those allocations deliver productive work. Inside allocated capacity, four runtime patterns commonly widen the gap between the allocated, active, and productive tiers, and separating them requires correlating device telemetry with execution timelines.

A zombie job retains an allocation after its useful execution has stopped, so its devices stay in the allocated tier while contributing nothing to the productive one, and they may or may not register as active. In a distributed run training a 175B parameter model, a rank failure can leave a 128-GPU group idle if worker cleanup or failure propagation does not complete. The surviving ranks may block inside a collective that will never finish, or a process left behind by a driver-level fault may resist reaping by the container runtime, so the devices stay allocated until a timeout or node-level remediation clears them. Persistent memory allocation combined with low SM activity is an investigation signal, not proof of a driver fault; valid input or communication waits can look similar. The operator must compare runtime health, step progress, and the workload's timeout contract, then verify that old workers have stopped before releasing or reassigning their devices.

The data starvation pattern lowers the active tier. It manifests as a high-frequency sawtooth wave in GPU utilization, where the device oscillates between 100 percent load (forward/backward pass) and 0 percent load (waiting for the next batch). While the average utilization might appear acceptable at 60 to 70 percent, the GPU is effectively idling for a third of its life. This occurs when the CPU-bound data pipeline cannot sustain the throughput required by the GPU, a common scenario in computer vision where JPEG decoding and augmentation saturate host CPUs. The data stall ratio of equation 4.4 quantifies the loss as the fraction of each step spent waiting for input. A timeline that aligns the idle intervals with input waits supports the diagnosis; utilization variance alone does not identify the cause. The remedy depends on whether decoding, storage, or transfer limits the input path.

A communication bottleneck can leave ranks waiting for collectives, which lowers the active tier, though communication kernels can also keep SM activity high and mask the wait in that tier. Per-rank gradient traffic depends on gradient precision, sharding, and the parallelism layout, so the ratio of step time to compute-only time (forward plus backward passes in isolation) is a quicker first bound on exposed overhead. If total iteration time exceeds 1.5× compute-only time, non-compute overhead exceeds 0.5× compute time, which is material but not dominant. Setting $T_{\text{sync}}(N)$ and T_{overlap} to zero in the Fleet Law (equation 1.3) leaves $T_{\text{step}}(N) = T_{\text{compute}}/N + T_{\text{comm}}(N)$, under which communication exceeds per-device computation only when the ratio is greater than 2×. Real steps also contain synchronization waits and input stalls, so profiling must separate communication from those before the ratio is read as a communication diagnosis. Placement can reduce exposed communication when it improves the relevant paths. Gradient accumulation can reduce synchronization frequency only when the runtime defers gradient exchange across accumulation steps.

On the serving side, memory fragmentation can prevent a replica from admitting new requests even when its HBM occupancy suggests free capacity, which lowers the productive tier without necessarily showing in the active one. Its signal is a replica that holds fewer concurrent requests than its occupancy predicts. The KV-cache allocator, rather than the aggregate occupancy percentage, decides whether a new sequence fits in contiguous or paged memory blocks. For the fleet scheduler, the allocator matters through its result, the number of concurrent requests a replica can hold, because that measured capacity sets how many replicas the serving envelope needs.

The debugging example connects scheduling mechanics to fleet governance. Technical policies (GPU type matching, gang scheduling timeouts, preemption grace periods) interact with organiza-

↪ Paged KV-cache allocation and its effect on replica concurrency are developed in section 10.7.3.

tional policies (team quotas, priority hierarchies, cost allocation) and human behavior (job template reuse, resource request habits, queue submission patterns). Effective fleet orchestration requires attending to all three layers simultaneously, because the most expensive scheduling bug is often not a software defect but a policy that was correct for the old infrastructure and was never updated for the new one.

8.13 Fallacies and Pitfalls

Those policy-infrastructure mismatches recur as a small set of misconceptions that can make an apparently healthy scheduler waste expensive fleet capacity. A cluster running at 98 percent utilization can look healthy while researchers have stopped submitting jobs because wait times have stretched into weeks. Each misconception below follows the same pattern, in which optimizing one visible metric degrades the outcome the fleet exists to deliver.

Fallacy: *More sophisticated scheduling algorithms always improve utilization.*

Engineers facing low utilization often reach for more advanced schedulers, assuming the algorithm is the bottleneck. As the debugging example in section 8.12 demonstrates, the root cause is frequently policy misconfiguration, not algorithmic limitation. A simple FIFO scheduler with correct policies (capability-based matching, backfill with timeouts, appropriate gang scheduling constraints) often outperforms a sophisticated scheduler with incorrect policies. Before upgrading the scheduler, audit the policies: confirm that users are not requesting resources they do not need, that heterogeneous resources are properly exposed, and that gang-scheduling timeouts are configured. The audit recurs as workloads evolve, because a preemption grace period that suffices for one model can be too short for another to save state, and placement weights suited to an earlier topology can strand capacity after an expansion.

Pitfall: *Treating all GPU-hours as equal when measuring utilization.*

An unlabeled "GPU utilization" percentage can conflate allocation with activity or useful work. For example, 80 percent of fleet capacity might be assigned to jobs while 15 percent of the fleet sits in groups still waiting for their remaining members and another 5 percent waits for input. Only 65 percent is then allocated to complete groups, and at most 60 percent executes useful training in that snapshot. Such a breakdown comes from scheduler and runtime evidence, which an activity counter cannot supply. The allocated, active, and productive tiers of section 8.12.1 keep these measurements distinct, with productive capacity tied to a stated workload objective and the job's fault-free rate.

Fallacy: *Gang scheduling is always necessary for distributed training.*

Gang scheduling prevents partial-allocation deadlock for a rigid synchronous group, but the required group need not equal the job's maximum requested size. Asynchronous workers may progress independently, while elastic training can begin with its minimum viable gang and resize at supported boundaries. Independent hyperparameter trials can also run separately. The scheduler gains flexibility by admitting the actual unit of progress, while preserving atomic allocation wherever synchronous dependencies require it. Changes to worker count or synchronization policy still require convergence validation.

Pitfall: *Setting static quotas based on peak demand.*

Organizations commonly set team quotas to handle worst-case demand, reasoning that teams need guaranteed access during crunch periods. If Team A's peak demand is 500 GPUs and Team B's is 300 GPUs, the cluster needs 800 GPUs with static quotas. In practice, peaks rarely coincide. A 600-GPU cluster can accommodate either peak when the other team's simultaneous demand leaves enough capacity, but it cannot guarantee both peaks concurrently. Hierarchical fair-share with borrowing uses that statistical opportunity only when the declared guarantees fit physical capacity and borrowed work can yield within the reclamation contract. In this example, sizing static quotas to the sum of the peaks provisions 800 GPUs where a 600-GPU shared cluster with borrowing covers either peak. The extra 200, a quarter of the larger cluster, are guarantees that sit unused whenever the peaks do not coincide.

Fallacy: *Spot instances are always cheaper for training.*

Actual savings depend on the billed overhead each interruption adds, not on the scenario's 60 to 90 percent discount band. Consider 15-minute checkpoint overhead to slow storage and 30-minute restart overhead to reacquire capacity and reload, far longer than a restart on a cluster's own file system, totaling 45 minutes of billed nonproductive time per event. If each event follows 2 hours of retained training progress, the nominal 65 percent discount becomes about 52 percent savings under the fixed-allocation model. This illustration assumes a successful save and no replay; a real interruption budget must also include discarded progress when the save cannot finish. The scheduler must separately account for unbilled capacity waits when evaluating deadlines. Sufficient billed overhead can erase the discount, regardless of the advertised rate.

Pitfall: *Ignoring topology when scheduling distributed training.*

Schedulers that ignore connectivity can place communication-intensive groups across paths with less available bandwidth or more contention. On a fleet with NVLink confined to each node, spreading a tensor-parallel group across nodes moves its traffic onto the inter-node fabric; other hardware can expose larger NVLink domains. Crossing a spine switch does not itself prove congestion, so the placement decision must compare effective path capacity and competing traffic. A measured throughput penalty accumulates over a long run and may exceed the delay of waiting for a more suitable allocation. The scheduler must weigh that gain against queue delay and fragmentation.

Pitfall: *Using elastic training to avoid gang-scheduling decisions.*

Elastic training complements gang scheduling by moving between viable execution groups. With a fixed 8-way tensor-parallel layout, each group still requires eight workers; changing the number of data-parallel replicas does not relax that requirement. A runtime that changes the tensor or pipeline layout must also redistribute state and establish a new viable group. The scheduler must honor the supported resize granularity, atomic admission requirements, and placement budget for each configuration. Elasticity adds feasible choices; it does not make an incomplete group productive or an arbitrary placement efficient.

Pitfall: *Treating high utilization as proof of cluster health.*

A high allocated-GPU percentage does not establish that the cluster has enough headroom to admit new work promptly, because waiting time climbs steeply as offered load nears capacity (figure 8.16). That queueing model does not turn a GPU-busy percentage into a prediction for a gang-scheduled fleet. Real headroom targets depend on arrival bursts, job sizes, placement constraints, and the queue-delay or serving-latency objective. Measuring those outcomes instead of applying a universal utilization threshold keeps platform teams from optimizations that improve a dashboard metric while delaying the work the fleet exists to complete.

8.14 Summary

Fleet orchestration turns data center capacity into useful ML work by deciding which jobs hold which accelerators, where, and for how long. Orchestration spends Coordination, the third axis of C^3, so that Compute and Communication capacity turn into useful progress. In Fleet Law terms (equation 1.3), placement sets the communication term $T_{\text{comm}}(N)$ of every step, and preemption and recovery, amortized over the run, enter the coordination term $T_{\text{sync}}(N)$. Queue wait and gang admission are coordination costs too, but they delay the first step rather than lengthening any step, so they appear in end-to-end completion time rather than in the per-step Fleet Law. Multi-tenancy and quota policy decide how much of that time allocated compute spends doing useful work. Hardware capability sets the feasible choices; the scheduler turns those choices into useful allocations by preserving ownership, coordinating recovery, and matching placement to workload demand.

Feasibility is the first filter. Gang size, topology locality, multi-dimensional fit, and accelerator compatibility decide which allocations a job can actually run, and only then can a fairness policy choose among them (principle 15). Quota borrowing extends the same reasoning across teams, since a guarantee is worth what the platform can reclaim within its response-time objective rather than the number printed on the quota.

Placement changes the physical paths available to a workload. Aligning model parallelism with the network hierarchy reduces exposed communication when placement is the limiting constraint,

but it cannot create fabric capacity that is not there. Elastic resizing and spot capacity widen the set of feasible choices by trading worker continuity for flexibility and price. ML-aware schedulers sharpen the choice among them with workload signals such as attained service, iteration boundaries, finish-time ratios, and goodput. Every one of these levers pays a transition or integration cost that must be measured against the progress it buys.

At serving time, the fleet scheduler preserves a capacity envelope of indivisible replicas with SLO headroom and moves the training-serving boundary on forecast demand, holding ready headroom for what the forecast misses. Request-level routing and batching stay with the serving system. When that envelope, the quotas around it, or the placements inside it go wrong, the visible symptom is usually a full queue beside idle accelerators. In practice, the root cause is more often a policy written for an earlier fleet than a defect in the scheduling algorithm.

Key Takeaways: Scheduling is systems engineering

- **Distributed scheduling is hard**: Cluster scheduling must coordinate allocation ownership across machines despite partial failures, network partitions, and stale observations. Control planes balance prompt recovery against the risk of authorizing a replacement while an old worker still runs.
- **Gang scheduling prevents deadlock but reduces flexibility**: Gang scheduling gives rigid execution groups atomic resource allocation to prevent hold-and-wait deadlocks, while backfilling and elastic resizing recover flexibility within the job's minimum viable gang and transition constraints.
- **Fairness is measured over feasible bundles**: DRF equalizes each tenant's dominant normalized share across GPUs, CPU, memory, and network, but only among bundles that already pass gang, topology, and compatibility constraints (principle 15). Slowdown and maximum feasible wait test whether the resulting policy avoids delay and starvation, and a quota guarantee holds only as far as lent capacity can be reclaimed in time.
- **Topology determines performance**: Where GPUs are placed within the cluster hierarchy changes the communication paths available to a job. Locality can reduce exposed communication, but waiting for a compact placement also costs time; compare end-to-end completion rather than iteration throughput alone.
- **Elastic and interruptible capacity must repay its transitions**: A forced preemption or spot interruption discards work since the last durable checkpoint and then pays reload and warmup, while a resize pauses the job. A resize pays only if the new scale persists long enough to amortize that pause, and spot capacity saves money only while the spot price, scaled by the ratio of billed to retained productive time, stays below the on-demand price.
- **Scheduling signals must match the objective**: Workload characteristics support different improvements: attained service addresses runtime uncertainty, iteration boundaries reduce interruption cost, finish-time ratios expose unfair delay, and goodput guides adaptive allocation. Each policy requires validation against the fleet's workload mix and recovery costs.
- **Useful utilization creates economic capacity**: Reducing avoidable idle time, often by updating policies written for an earlier fleet rather than by changing the scheduling algorithm, can deliver more useful work from an existing fleet. Distinguish allocated, active, and productive capacity before valuing an improvement, and preserve the headroom required for recovery and latency targets.

Placing a tensor-parallel gang makes the displacement of overhead (principle 14) concrete. Colocating the gang keeps its collectives on NVLink and shrinks the exposed communication of every step, but the scheduler pays for that locality in coordination time while it waits for a feasible contiguous allocation. The bottleneck moves from collective bandwidth to queue delay or fragmentation

rather than disappearing, so the right placement is the one that minimizes end-to-end completion time, not iteration time alone.

> **What's Next: From allocation to useful work**
>
> An allocation shows that a job holds capacity, not that the capacity produces useful work. A training job can hold a complete allocation while it restarts, waits for input, or exposes communication in every step. A serving replica can hold a full node while its decode steps stream weights and leave most of its arithmetic idle. The productive tier says how much of an allocation became useful work but not which part of a job's life lost the rest. The owners of scheduling, recovery, and execution therefore each see their own loss, and none can see which is largest. Those separate losses become one product whose factors can be compared once every hour of capacity is followed from the moment the fleet offers it until it becomes useful work, with each loss charged to the boundary where it occurs.

↪ How much of an allocation becomes useful work, and whether scheduling, runtime, or the program loses the rest, is developed in chapter 9.

III

THE FLEET IN PRODUCTION

Part III

8

Fleet-in-Production Principles

Parts I and II built the fleet, coordinated its distributed state, and executed the training run. Training, however, opens a model's lifecycle rather than completing it, and production deployment is where the model must repay its cost under real demand. Part III shifts the focus from the single, large training run orchestrated across a dedicated cluster to the deployed fleet. That fleet comprises the thousands of geographically distributed servers that answer requests in real time, the billions of edge devices that run and adapt models locally, and the operational machinery that keeps every deployed version accountable. The transition changes the engineering constraints. A training run must finish a coordinated workload. A deployed fleet must sustain useful throughput on a dynamic request stream while meeting latency targets, keep models current as data and devices change, and govern how each change reaches production.

A serving fleet depends on the runtime behavior of each replica (one deployed copy of the model, potentially spread across devices), whose execution time and memory footprint determine how many requests it can accept and how much capacity the fleet must provision. Part III therefore begins with fleet efficiency, which follows the fleet's capacity through scheduling, runtime, and execution to measure how much of it becomes useful work. Serving at scale uses that measured capacity to coordinate request queues, batching, serving precision, and distributed model state. Edge intelligence extends the argument to fleets of varied client devices with intermittent connectivity, including federated learning that coordinates training while data remains on participating devices. Production operations then governs rollout, monitoring, and recovery across the deployed fleet.

The principles below preview the constraints that each chapter of Part III makes quantitative, turning replica-level mechanisms into fleet-wide capacity and latency models rather than assuming that local efficiency alone guarantees an efficient fleet. Before optimizing any serving mechanism, the part establishes a measure of how much of a fleet's capacity becomes useful work. Without an end-to-end accounting rule, losses in cluster scheduling, runtime recovery, and kernel execution hide behind separate team metrics.

Principle 16: The ML goodput decomposition

Invariant: A fleet's ML productivity goodput, the fraction of its capacity-time that becomes useful work, is the product of three factors, each counted at a stated boundary so that every loss belongs to exactly one. Scheduling goodput $\eta_{\text{scheduling}}$ is the share of capacity held by complete allocations. Runtime goodput η_{runtime} is the share of allocated time retained as progress after startup, input stalls, checkpoint pauses, restarts, and rework. Program goodput η_{program} is the fraction of the ideal rate that retained steps achieve, measured as MFU.

$$\text{MPG} = \eta_{\text{scheduling}} \times \eta_{\text{runtime}} \times \eta_{\text{program}}$$

Implication: Lost capacity must be located before it is recovered, because each factor has a different owner. A faster kernel recovers nothing that a job loses waiting for a complete gang, and a better scheduler recovers nothing that a job loses running below its hardware's rate. A change that moves a loss from one factor to another helps only if the product rises. The decomposition applies to a training step and to a serving replica alike; serving adds the latency targets that each completed request must meet.

Over a deployed model's lifetime, serving can consume more of that capacity than training did, and serving also carries the book's scaling thesis into the request path. When quality is bought with test-time compute, by spending more computation on each query, the scaling curve that drove the training fleet moves into serving, and its cost recurs with every request.

> **Principle 17: Serving cost dominance**
>
> **Invariant**: For high-traffic production models, cumulative inference cost can exceed the one-time training cost over the model's lifetime.
>
> **Implication**: Inference efficiency often becomes the dominant lifecycle optimization target. When traffic volume, product lifetime, and retraining cadence make serving cost dominate, engineers should focus quantization, distillation, and sparsity on the inference path. The quantity to reduce is the cost per completed request at a fixed quality and latency target, not the cost of an individual operation.

Within each request, the mechanics of generation impose the first limit. Autoregressive generation produces each new token from the preceding tokens while a key-value cache retains intermediate attention state for reuse, so a single serving replica can saturate either memory bandwidth or memory capacity long before its arithmetic units are full.

> **Principle 18: The autoregressive bottleneck**
>
> **Invariant**: In low-batch dense transformer decode, weight streaming often makes throughput memory-bandwidth bound because every decode step must stream the model weights from device memory to the compute units. Higher batch sizes amortize weight streaming, and key-value (KV) cache traffic shifts the binding constraint at long context.
>
> **Implication**: Throughput initially improves with batch size by amortizing weight reads across requests, until per-device latency budgets, KV cache capacity, or sustained compute become the binding ceiling. Serving systems therefore need batching, cache-management, and memory-layout policies that keep bandwidth useful without violating latency SLOs.

At scale, routing policy becomes another lever for protecting latency. Even well-tuned replicas develop long queues if incoming queries land on busy servers while neighboring replicas have spare capacity.

> **Principle 19: The power of two choices (P2C)**
>
> **Invariant**: In the idealized process that sequentially places R independent requests into R replicas, with each request sampling replicas independently and uniformly at random, querying two replicas and selecting the least-loaded one reduces the maximum load from the one-choice bound $\Theta(\log R / \log \log R)$ to $\Theta(\log \log R)$ with high probability.
>
> **Implication**: Simple randomized load-balancing algorithms produce dramatically lower maximum-load bounds with $\mathcal{O}(1)$ coordination overhead. Under serving queueing assumptions, the lower maximum load translates into reduced tail latency at scale. Serving requests differ widely in remaining work and load probes go stale, so the two probes must read a work signal, such as active tokens or KV-cache pressure, rather than a request count.

Moving learning onto phones and other battery-powered devices adds a constraint that no data-center scheduler faces. The device belongs to its user, so background learning runs only when the device can spare it.

Principle 20: The participation selection principle

Invariant: An edge device contributes to on-device or federated learning only in windows its admission policy accepts, when the predicted update energy fits the remaining energy slack, the added power stays within the thermal limit, and foreground latency still meets its deadline (equation 11.2). Those conditions correlate with hardware class, charging habits, and connectivity, so the population a shared model learns from is the admitted population, not the deployed one.

Implication: The admission policy that protects each device also selects the training population, which makes admission a fairness decision as well as a resource decision. A fleet must measure participation and completion rates per device and user cohort against the deployed population and correct the skew through stratified selection, participation quotas, or reweighting. Loosening admission is not a remedy, because it spends the energy, thermal, and latency budgets the policy exists to protect.

Whether a model is trained in the data center or across a population of devices, every change must eventually be released to the deployed fleet, and at fleet scale that release needs a single standard of evidence. Shipping a candidate on scattered benchmarks or informal spot checks leaves operators unable to tell whether a production incident reflects a bad release or an unmeasured shift in traffic.

Principle 21: The evaluation contract principle

Invariant: At fleet scale, a release decision is valid only when all of its evidence refers to one versioned evaluation contract that pins the model and its dependencies, the workload, the metrics and statistical design, the acceptance thresholds, the rollback conditions, and the accountable owners. Reliability under that contract is a vector of budgets over availability, latency, freshness, and quality. With B_i denoting the fraction of observations that violate target i divided by the violating fraction that target allows, the service is inside its reliability envelope only when

$$\max_i B_i \leq 1.$$

Implication: No single accuracy number can authorize a release. Averaging the dimensions would hide a quality regression behind perfect uptime, and evidence gathered on different datasets, load profiles, or configurations certifies no single candidate. Thresholds and rollback conditions are fixed before exposure, an exhausted budget in any dimension stops promotion, and a change that leaves the contract's validated envelope reopens the decision.

Part III takes the trained model from the cluster to the world. The path begins with fleet efficiency, which charges every lost hour of capacity to scheduling, runtime, or the program and so determines how much of the fleet's capacity becomes useful work. It then moves to serving at scale, where request batching, cache management, serving precision, routing, and sharding turn single-device efficiency into fleet-level latency control. Edge intelligence extends this discipline to phone-class devices that adapt models locally and coordinate through federated learning under unequal memory, power, and connectivity budgets. Operations at scale closes the loop by monitoring the deployed fleet, coordinating platform ownership, and responding when production behavior diverges from design assumptions. Measured against the Fleet Law, these chapters shift from the training step to the serving request. Fleet efficiency charges each term of a step to the factor that loses it and shrinks the exposed ones. Serving at scale trades batching against the queueing and routing delays that play the role of $T_{\text{sync}}(N)$. Edge intelligence faces communication costs set by intermittent wireless links rather than a data-center fabric. Operations bounds the coordination cost of every rollout and rollback. Together, these chapters develop the operational discipline that turns a trained model into a durable production system.

9

Fleet Efficiency

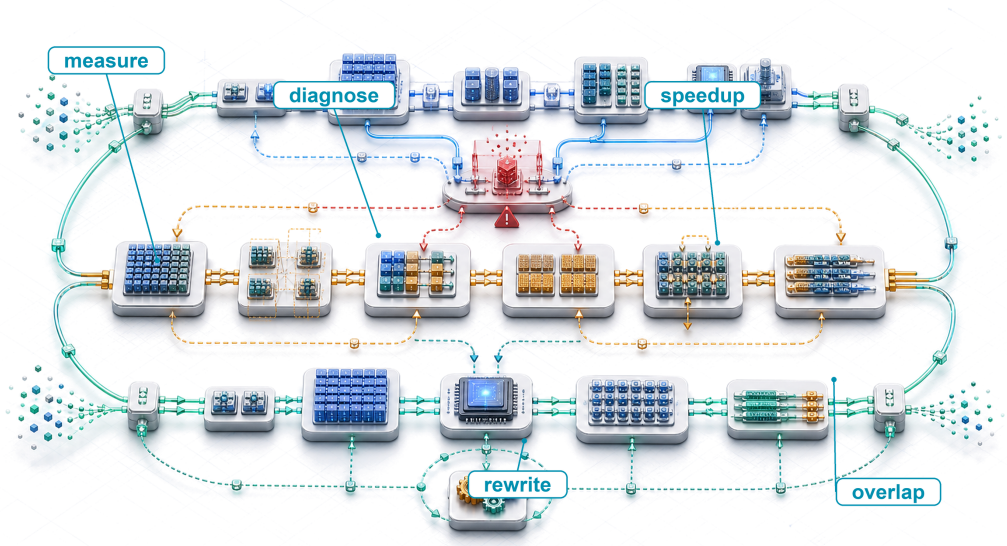

9.1 Where Fleet Capacity Goes
9.2 ML Goodput
9.3 Measuring Goodput
9.4 Scheduling Goodput
9.5 Runtime Goodput
9.6 Program Goodput
9.7 Recovering Program Goodput
9.8 The Whole Fleet
9.9 Trading Between Factors
9.10 Fallacies and Pitfalls
9.11 Summary

Purpose

Why can a fleet that is fully allocated and visibly busy still turn only a fraction of its compute into training progress and served requests?

Expensive accelerators can spend much of a training step or a request waiting for data, software dispatch, or communication while their arithmetic units remain underused. Adding more machines may raise cost without removing the delay, and a faster isolated kernel may leave the step's or request's critical path unchanged. The engineering stakes are useful fleet capacity, the training progress and completed service an existing fleet produces while preserving model quality. A dashboard that counts every held accelerator as busy cannot see that waste, so a fleet can look fully used while turning only a fraction of its capacity into useful work. These interactions make optimization a question of measured constraints rather than a search for the fastest component. On one machine, a local improvement can be tested against a relatively contained execution path. Across a fleet, its value also depends on scheduling, recovery, communication, synchronization, and the distribution of work across devices. The cost of a mistaken diagnosis grows when the same inefficient execution repeats across many replicas and many weeks of training, and when the remedy is applied by a team that does not own the loss. A saving that looks large in isolation is worth only the part that survives along the complete step or request path. In C^3 terms, fleet efficiency holds every compute gain accountable to the communication and coordination costs that decide how much of a fleet's capacity becomes useful work.

Learning Objectives

- Explain how a job's or fleet's useful output decomposes into scheduling, runtime, and program goodput, with each loss charged to exactly one factor
- Measure goodput under a reproducible contract and localize the lost factor from scheduler records, job timelines, and step traces
- Attribute scheduling and runtime losses using the owners' results, without re-deriving them
- Measure program goodput with MFU and MBU and find the Fleet Law term a fix must shorten
- Select fleet-level program levers (compiler partitioning, overlap under contention, compilation) and recognize kernel-level levers by the term they shorten
- Evaluate a change that trades one factor against another by its effect on the product, and place the result on the efficiency frontier

9.1 Where Fleet Capacity Goes

A fleet can be fully allocated and visibly busy while most of its capacity never becomes training progress or completed service. When thousands of accelerators run a single workload, losses accumulate before a job starts, between the steps it executes, and inside the steps it keeps. The size of that gap is clearest in a run whose engineering is not in doubt.

Lighthouse 9.1: Archetype A (GPT-4/Llama-3): Held capacity versus progress

Meta's pretraining of Llama 3 405B on 16,384 H100 GPUs is one of the best-documented large training runs, and its two most quoted efficiency numbers measure different things. Across a 54-day snapshot, the job kept more than 90 percent of its elapsed time as effective training time (time spent on useful training) despite hundreds of unexpected interruptions, as section 7.9.4 records. In the steps it kept, its arithmetic achieved a model FLOPs utilization (MFU) of about 41 percent of the accelerators' peak rate (Dubey et al. 2024). Multiplied, the two figures show the run turning roughly 37 percent of the peak-rate capacity it held into training progress. Because the two figures come from different measurement windows, the product is only an estimate. Yet it shows one of the most carefully engineered runs on record losing about two-thirds of the capacity it held, mostly to steps that ran below the hardware's rate rather than to the failures that dominate its reputation. Its two figures therefore illustrate what a well-run job of this archetype loses during execution.

Even that one-third figure overstates what a shared cluster delivers, because the Llama 3 numbers begin only after the job holds its accelerators. On a multi-tenant fleet, the setting of chapter 8, capacity is also lost before a job starts. Accelerators sit unallocated when no queued job fits the free pool, or a job holds part of its gang (the set of accelerators that must all start together) while waiting for the remaining members and cannot execute a single step. The orchestrator's accounting separates allocated capacity from productive work, yet it cannot see where inside a running job the missing work went. Allocated capacity is where fleet-efficiency accounting begins, not evidence that the fleet used those accelerators well.

↩ The allocated, active, and productive tiers of fleet resource accounting are established in section 8.12.1.

Earlier chapters examined the mechanisms of a fleet one layer at a time: data delivery (chapter 4), collective communication (chapter 5), parallel work partitioning (chapter 6), fault recovery (chapter 7), and cluster scheduling (chapter 8). Each subsystem reports performance in its own local unit—input stall ratio, link bandwidth, scaling efficiency, availability, or cluster occupancy—and none of those isolated metrics tells an organization what fraction of its fleet investment becomes useful progress. Worse, when subsystem reports are read separately, their denominators overlap in some places and leave gaps in others. An hour lost to a restart can be penalized in two reports at once, while accelerators held idle by an incomplete gang vanish between the scheduler's log and the job's step timer. Charging each lost hour exactly once requires a single end-to-end accounting framework that

follows every accelerator-hour from offered hardware to useful output and ranks every optimization lever across the stack on one common scale. That framework opens the Serving band of figure 1.4 by asking how much of the infrastructure built in the lower bands becomes useful work.

To make that accounting concrete, one recurring training job carries the numbers through this chapter: pretraining Llama 2 70B, a smaller Archetype A model than the archetype's 175B reference, on 128 GPUs (16 nodes of eight GPUs each) in a shared cluster. The job uses 8-way tensor parallelism within each node and data parallelism across nodes. Because it competes with other teams for complete gang allocations, absorbs cluster failures and preemptions, and executes steps whose exposed communication grows with node count, every stage of fleet loss appears in a single workload on a single pool. The job is a constructed illustration rather than a measured run, so its numbers are modeled from stated assumptions.

↩ The Fleet Law, where these multiplicative efficiency factors are first previewed, is established in section 1.5.4.

A serving replica of the same model provides the running example's inference twin. The replica spans one node of eight H100 GPUs, so a single tensor-parallel node forms one complete serving unit. Because a dense Archetype A transformer cannot fit on one accelerator, the eight GPUs exchange activations in every layer. The cluster scheduler must place the replica as an indivisible gang, the replica must load weights and warm up before taking traffic, and its decode steps can leave most of the node's arithmetic idle while weights stream from high-bandwidth memory (HBM). Its capacity losses therefore arise at the same stages as those of the training job.

↩ Archetype A and its 175B-parameter reference model are established in section 1.6.1.1.

The serving twin also marks where fleet-efficiency accounting stops and request-level serving design begins. In production serving, requests share execution across dynamically formed batches, occupy key-value cache memory, and must satisfy tail-latency targets; chapter 10 develops how a serving engine manages uneven arrivals under those constraints. Here, the boundary test is whether a loss would persist even under a steady, uncontended stream of work. A replica held off-line while loading weights, or a decode step that reads weights more slowly than the node's memory bandwidth allows, fails to use held hardware even without arrival bursts and is counted here over windows of hours to weeks. By contrast, a request delayed behind an oversized batch is a queueing phenomenon caused by uneven arrivals and belongs to serving analysis. Headroom that a replica holds in reserve to absorb traffic spikes appears in goodput accounting as ready capacity running below its ideal rate, whereas choosing how much headroom a latency target requires is a serving decision.

This chapter builds the measurement and optimization discipline that turns those scattered losses into one engineering budget. We begin by defining *ML goodput* (section 9.2), which factors useful fleet output into scheduling, runtime, and program goodput on non-overlapping boundaries. We then establish the reproducible measurement contract and hierarchical fleet profiling needed to localize which factor moved without drowning the cluster in trace data (section 9.3). From there, we trace the three factors in chronological order. We attribute scheduling losses (section 9.4) and runtime losses (section 9.5) to the orchestration and fault-tolerance mechanisms developed in earlier chapters, and then diagnose (section 9.6) and recover (section 9.7) the program-goodput losses that this chapter owns directly. Finally, we compose all three factors into a whole-fleet waterfall and cost per useful accelerator-hour (section 9.8) and evaluate optimizations that trade one goodput factor against another along the fleet's efficiency frontier (section 9.9).

9.2 ML Goodput

Across a fleet, capacity disappears at three successive stages in the life of a job. First, a job must acquire a complete set of devices: an accelerator produces no progress while it sits unallocated, while cluster fragmentation strands it in a shape no queued gang can use, or while a partially allocated gang holds it idle waiting for the rest of its peers. Second, once a job holds a complete allocation, its execution must produce durable progress: initialization, input stalls, synchronous checkpoint pauses, failure recovery, and steps re-executed after a rollback all consume held device-hours without advancing the last saved checkpoint. Third, the steps that do survive in a checkpoint must execute near the rate the hardware can sustain. A step that exposes collective communication, waits on a straggling rank or a pipeline bubble, stalls on memory traffic, or dispatches long chains of tiny kernels occupies healthy accelerators while extracting only a fraction of their capability.

9.2.1 The three goodput factors

Run-level goodput, introduced with the Coordination dimension in section 1.5.1.3, tracks a single job across its execution and discounts time lost to failures and replayed steps. Because that measure starts only after the job holds its accelerators and stops at the boundary of a completed step, it misses both the capacity the job never received and the execution efficiency of the steps the job kept. Google's fleet-efficiency framework extends the measurement window backward to the capacity offered by the fleet and forward to the ideal execution time of each retained step (the time its required work would take at the peak rate of the hardware resource that limits it), factoring useful fleet output into three multiplicative terms (Wongpanich et al. 2025).[1]

[1] **ML productivity goodput origin**: Google introduced the three-factor decomposition to measure its Tensor Processing Unit (TPU) fleet, calling it an iron law of performance for ML fleets by analogy with the processor iron law that splits execution time into instructions, cycles per instruction, and time per cycle (Emer and Clark 1984; Wongpanich et al. 2025). The paper reports the factors across the whole fleet and segmented by job size and by workload phase, with real-time serving and bulk inference alongside training. The boundaries themselves name no accelerator, so this chapter applies them to any accelerator fleet, for serving replicas as well as training jobs.

> Definition 9.1: ML productivity goodput (MPG)
>
> ***ML productivity goodput (MPG)*** is the fraction of a fleet's capacity-time that becomes useful work at the hardware's ideal rate, expressed as the product of scheduling, runtime, and program goodput, each measured on a fixed accounting boundary so that every lost interval is charged to exactly one factor (Wongpanich et al. 2025).
>
> 1. **Significance**: It converts the losses that scheduling, recovery, and execution each report in their own units into one fraction of fleet capacity, so a loss at one boundary can be compared directly with an equal loss at another, and an accelerator-hour's price divided by it is the cost of one ideal-rate hour of useful work.
> 2. **Distinction**: Unlike utilization, which counts an accelerator as used whenever it is allocated or busy, it counts only retained work at its ideal rate. Unlike run-level goodput, it also charges the capacity a job never received and the rate at which its kept steps ran.
> 3. **Common pitfall**: A frequent mistake is to multiply factors measured over different windows or against overlapping denominators, such as a run-averaged MFU that already includes restart time and a runtime factor that charges the same restarts again. The product is meaningful only when each factor divides by exactly what the previous one passed.

The three factors follow a job chronologically, and each factor takes as its denominator the capacity-time that survived the stage before it. **Scheduling goodput** ($\eta_{\text{scheduling}}$) is the fraction of fleet capacity-time held by complete allocations—intervals in which a job holds every resource its synchronous execution requires, which the fleet-efficiency literature calls all-allocated time (Wongpanich et al. 2025). **Runtime goodput** (η_{runtime}) is the fraction of that all-allocated time spent on steps that a checkpoint retains, excluding initialization, input stalls, synchronous checkpoint pauses, failure recovery, and work discarded after the last checkpoint; its numerator is productive time. **Program goodput** (η_{program}) is the ratio of the predicted ideal execution time of those retained steps to their actual productive time, capturing the efficiency of code, compiler, and communication against a named hardware ideal. In the Llama 3 405B run, effective training time serves as runtime goodput and MFU serves as program goodput; on a shared multi-tenant cluster, a scheduling factor below one multiplies both. Equation 9.1 combines the three factors into MPG:

$$\text{MPG} = \eta_{\text{scheduling}} \times \eta_{\text{runtime}} \times \eta_{\text{program}} \tag{9.1}$$

The three factors multiply cleanly because each factor's denominator is the preceding factor's numerator. Scheduling goodput divides held all-allocated time by offered capacity-time, runtime goodput divides retained productive step time by all-allocated time, and program goodput divides predicted ideal time by productive step time. Adjacent numerators and denominators cancel, leaving the ideal-rate time of useful work divided by the total capacity-time the fleet offered (figure 9.1). MPG is therefore a dimensionless fraction of fleet capacity rather than a raw throughput rate: useful fleet output equals peak capacity multiplied by MPG, and the effective cost per unit of useful work equals the cost of raw capacity divided by MPG. That telescoping cancellation also defines when the product is valid. If two factors include the same lost interval in their denominators, that loss is

counted twice; if a gap falls between one factor's numerator and the next factor's denominator, the loss goes uncounted in both.

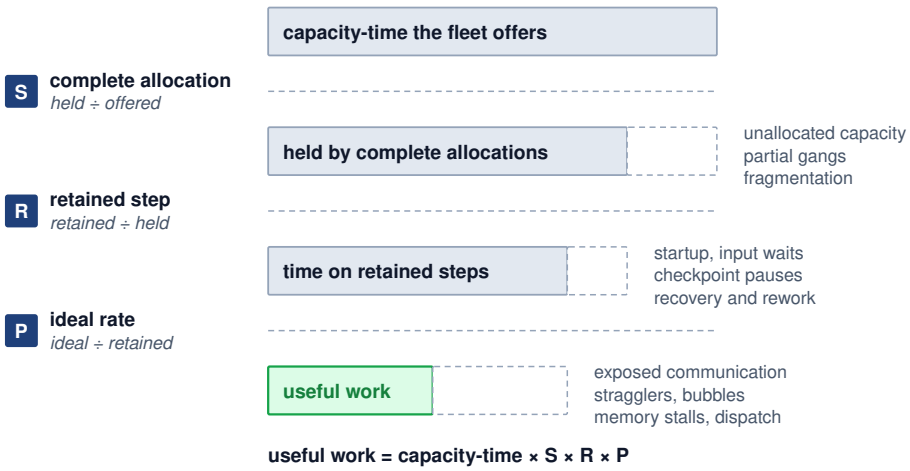

Figure 9.1: ML Goodput as a Chain of Three Boundaries: Capacity-time becomes useful work only by crossing three boundaries in the order a job meets them, and each factor divides what crossed its boundary by what reached it. The dashed slice at the end of each row is the capacity-time lost at the boundary above that row, and the losses beside it are charged there and at no other boundary, so the three factors multiply into the share of offered capacity-time that becomes useful work. Widths are schematic, not measured.

Each factor captures the losses from one stage of a job's lifecycle (section 9.1), and the three goodput boundaries cut across the C^3 taxonomy rather than mapping one-to-one onto its axes. Hiding collective communication behind computation raises program goodput, whereas checkpointing and failure recovery—which the Fleet Law places under Coordination—reduce runtime goodput. Goodput factors identify *when* in a job's execution capacity was lost, while the C^3 dimensions identify *which physical mechanism* caused the loss; diagnosing a fleet bottleneck therefore names the goodput factor first and the C^3 mechanism inside it second.

Because the three factors multiply, the share of total fleet capacity lost at any boundary equals that boundary's fractional loss weighted by the capacity share that reached it. Equation 9.2 partitions the total efficiency gap $1-\eta_{\text{scheduling}}\eta_{\text{runtime}}\eta_{\text{program}}$ into three additive shares—the scheduling loss on offered capacity, the runtime loss on held capacity, and the program loss on retained step time:

$$\begin{aligned} 1-\eta_{\text{scheduling}}\eta_{\text{runtime}}\eta_{\text{program}} = {} & (1-\eta_{\text{scheduling}}) \\ & +\eta_{\text{scheduling}}(1-\eta_{\text{runtime}}) \\ & +\eta_{\text{scheduling}}\eta_{\text{runtime}}(1-\eta_{\text{program}}) \end{aligned} \tag{9.2}$$

The three right-hand terms sum to the total gap, reflecting the rule that every lost device-hour is charged at exactly one boundary. Although the product $\eta_{\text{scheduling}}\eta_{\text{runtime}}\eta_{\text{program}}$ is commutative, the additive share attributed to each stage depends on the order in which capacity passes through the boundaries. The decomposition therefore follows the chronological order in which a workload encounters the three losses. Applied to the published Llama 3 405B telemetry, equation 9.2 shows immediately which boundary consumed most of the cluster's held capacity.

Napkin Math 9.1: Where Llama 3's held capacity went

Problem: Llama 3 405B trained on a dedicated cluster of 16,384 GPUs, kept more than 90 percent of its elapsed time as effective training time, and ran its kept steps at about 41 percent

MFU (Dubey et al. 2024). What share of the capacity it held was lost at the runtime boundary, what share at the program boundary, and what share became progress at the ideal rate?

Variables:

- Scheduling goodput $\eta_{\text{scheduling}}$ is one, because the cluster was dedicated to the run, so the first term of equation 9.2 vanishes.
- Runtime goodput η_{runtime} is about 0.90, from effective training time (a stated lower bound).
- Program goodput η_{program} is about 0.41, from MFU at full scale.

Math:

- Runtime boundary, $1-\eta_{\text{runtime}}$: $1-0.90=0.10$.
- Program boundary, $\eta_{\text{runtime}}(1-\eta_{\text{program}})$: $0.90\times 0.59\approx 0.53$.
- Useful work, $\eta_{\text{runtime}}\eta_{\text{program}}$: $0.90\times 0.41\approx 0.37$.

Result: About 10 percent of the held capacity was lost at the runtime boundary, about 53 percent at the program boundary, and about 37 percent became progress at the ideal rate.

Systems insight: Removing every interruption would have recovered at most a tenth of the capacity the run held, while the steps it kept left more than half unused, so for one of the most carefully engineered runs on record the program boundary, not reliability, held most of the gap.

This calculation treats Meta's reported effective training time as runtime goodput.[2] That identification holds when the job retained its full allocation throughout the elapsed window and when input stalls were already excluded from effective time. The calculation's second identification, using MFU as program goodput, holds when the MFU was measured over those same retained steps, whereas Meta's MFU comes from a different measurement window, so the notebook's shares are estimates.

9.2.2 Program goodput and its measure

Hardware activity counters cannot measure how efficiently a retained step ran, because an accelerator registers as busy whether it is waiting on a collective, recomputing dropped activations, or executing arithmetic far below its peak rate. Measuring program goodput instead requires comparing actual step time against a named hardware ideal for the resource that limits the workload. For training steps and for the prefill pass over a prompt, arithmetic dominates, so the ideal time of a step is the time the accelerators would need to execute the model's required mathematical operations at their peak rate. MFU measures against that ideal by counting only the operations the model architecture requires.[3]

In Google's fleet-efficiency methodology, the ideal operation count is fixed before compilation from the program as written—before the compiler fuses kernels or rematerializes activations—so that a compiler optimization that shortens the step raises program goodput instead of shifting its own baseline (Wongpanich et al. 2025). Program goodput is not measured against a per-operation roofline either, because a per-operation roofline rewards kernels that run near their local ceiling and would therefore penalize a rewrite that executes fewer operations at lower utilization while finishing the step sooner. The roofline remains a diagnostic for individual kernels (section 9.6.3) rather than the efficiency measure of a whole program.

Definition 9.2: Model FLOPs utilization (MFU)

Model FLOPs utilization (MFU) is the fraction of the hardware's theoretical peak throughput (R_{peak}) consumed by FLOPs that directly advance model training or inference, excluding overhead from recomputation, padding, and synchronization.

1. **Significance**: Measured over a job's retained steps against the arithmetic ideal, MFU is that job's program goodput. A one-percentage-point MFU increase across 10,000 identical

[2] **Effective training time**: Meta's report defines it as the time spent on useful training over the elapsed time and states only that it exceeded 90 percent over a 54-day snapshot (Dubey et al. 2024). It is a per-job figure, so it says nothing about capacity the job never held, and the report does not say whether input waits count as useful. Its MFU comes from the scaling-configuration table rather than from that snapshot, which is why the measurement contract (section 9.3.1) requires every factor to name its window.

[3] **MFU and HFU**: The PaLM report proposed MFU in place of hardware FLOPs utilization (HFU), which counts every executed operation, because the executed count depends on implementation and compiler choices, most visibly rematerialization, which recomputes activations to save memory (Chowdhery et al. 2022). MFU's reference rate depends only on the model architecture and the published peak rate, so runs of one model on different software stacks compare directly. For goodput accounting, the choice means that recomputation which lets a layout fit in memory is charged to program goodput as lost time rather than credited as work.

GPUs adds useful throughput equivalent to 100 GPUs at peak, and whether that saves money depends on whether the recovered capacity displaces hardware or shortens paid runtime.

2. **Distinction**: Unlike hardware utilization, which reports how often the accelerator is busy, MFU relates the model's required FLOP rate to hardware peak. It excludes recomputation and padding from useful work even when they make a configuration feasible or faster.
3. **Common pitfall**: A frequent misconception is that high GPU utilization implies high efficiency. A system can show 90 percent hardware utilization while achieving 30 percent MFU when busy intervals include communication or arithmetic far below peak. MFU exposes the gap but does not explain its cause, which traces must find.

The equality between MFU and program goodput under an arithmetic ideal follows directly from that definition. A retained step across N accelerators performs O_{useful} of required model work, which would take $O_{\text{useful}}/(NR_{\text{peak}})$ at peak throughput, whereas the step actually runs in wall-clock time T_{step}. Equation 9.3 expresses program goodput as the ratio of ideal step time to actual step time:

↩ The compute axis, where the memory wall limits MFU, is established in section 1.5.1.1.

$$\eta_{\text{program}} = \frac{O_{\text{useful}}/(NR_{\text{peak}})}{T_{\text{step}}} = \text{MFU} \tag{9.3}$$

That identity holds only when MFU is measured over the retained steps that survived the runtime boundary, with input stalls already excluded. Published MFU numbers are often averaged over an entire wall-clock run that includes restarts and synchronous checkpoint pauses; multiplying such a run-averaged MFU by a separate runtime goodput factor would penalize those interruptions twice. Hardware FLOPs utilization (HFU), which counts all executed floating-point operations including activation recomputation, remains a useful diagnostic alongside MFU because the gap between HFU and MFU isolates the cost of recomputation. In goodput accounting, that recomputation gap is a program goodput loss even when trading extra compute for memory capacity is the right engineering choice. Similarly, when capacity planners size a training run in accelerator-hours by scaling peak throughput by retained-step MFU (section 2.4.1), they must still multiply by runtime and scheduling goodput to account for the hours lost outside steady-state steps.

The MFU that published training runs report stays below half of peak across every generation of frontier hardware (figure 9.2). Even for the same model architecture, expanding Llama 3 405B onto twice as many GPUs lowers its reported MFU, which the report attributes to the smaller batch per data-parallel group needed to hold the global batch constant.

An arithmetic ideal does not fit workloads that the hardware cannot make compute bound. An autoregressive decode step generating one token for a small batch must stream every parameter of the model from HBM to perform only a handful of operations per weight. Memory bandwidth therefore sets the step's floor (section 2.3.1.3), and at that bandwidth floor the accelerator's matrix units run at well under 1 percent of peak arithmetic throughput. Charging that arithmetic shortfall to the program would blame kernel implementation for a gap that only a larger batch can close—and maximum batch size is constrained by the serving latency target rather than by the kernel. For bandwidth-bound decode, the natural reference for a single replica is the weight-only bandwidth floor: the time required to stream the model's weight bytes at peak memory bandwidth. Program goodput against that ideal is **model bandwidth utilization** (MBU), defined as the achieved rate of weight-byte transfers divided by peak memory bandwidth. The common MBU definition also counts key-value cache bytes in the floor. This chapter counts weights alone, so reading the key-value cache lengthens the measured step without entering the floor and counts against the replica as a program loss, which keeps cache growth visible. At long context, where cache reads can rival the weight bytes, a low MBU can therefore reflect cache size rather than kernel quality. A decode step with high MBU and low MFU has saturated the limiting hardware resource rather than wasted it.

Choosing a named ideal also determines what can be aggregated across a multi-workload cluster. A fleet that evaluates training against peak arithmetic and decode against peak memory bandwidth cannot average the two into a single program goodput factor, because the two denominators measure

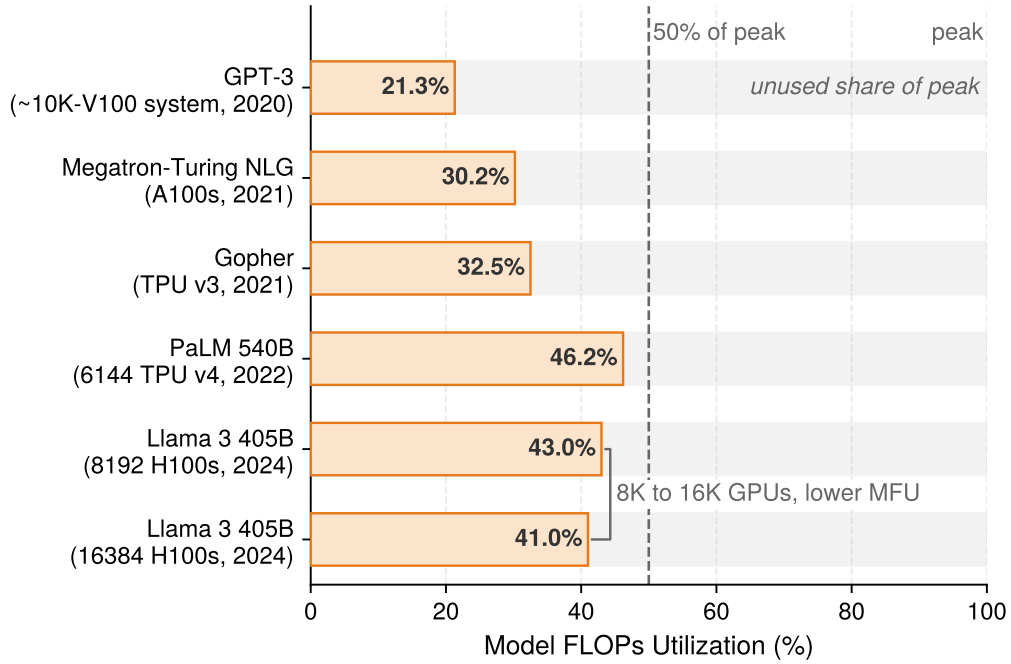

Figure 9.2: The MFU Progression: Model FLOPs Utilization across published training systems, ordered chronologically. Reported MFU rises from 21 percent (GPT-3, 2020) to 46 percent (PaLM, 2022), with later illustrated runs near 40–45 percent. These systems differ in hardware and workload, so the comparison does not isolate the effect of parallelism or overlap. The Llama 3 data points show lower reported MFU at the larger scale: MFU drops from 43 percent at 8,192 GPUs to 41 percent at 16,384 GPUs. Interpreting the difference requires the batch, execution configuration, and measurement convention. Data sources: Chowdhery et al. (2022), Dubey et al. (2024).

different physical resources. Fleet-wide reporting must therefore either keep program goodput segmented by workload class or evaluate every workload against the arithmetic ideal and recognize that small-batch decode carries a structural batch-size penalty. Google's fleet-efficiency framework takes the second path for its fleet-wide aggregate, measuring all workloads against a compute-based ideal (Wongpanich et al. 2025), while MBU provides a direct diagnostic for individual bandwidth-bound decode replicas. Throughout this chapter, MFU serves as the default program goodput metric because it composes across the fleet, and MBU is reported alongside it whenever memory bandwidth sets the hardware ceiling.

For the chapter's running Llama 2 70B training job, program goodput is the fleet MFU of its retained steps across all 128 GPUs. In the constructed example, that fleet MFU is about 48 percent, down from the 65 percent of a single-node baseline, because distributing each step across 16 nodes introduces inter-node collective communication that a single node never pays. Both MFUs rest on illustrative step times for a microbatch-scale step of about 2,215 tokens per data-parallel replica, so one optimizer update accumulates many such steps. The program goodput of 0.48 is therefore a per-microbatch value. The data-axis gradient reduction runs once per optimizer step and is amortized across the accumulated microbatches (section 9.7.2), which lowers program goodput over a full optimizer step to about 0.456. This chapter carries the per-microbatch value forward and states the full-step value wherever the job's factors are composed. Section 9.6.2 isolates and quantifies the scaling loss.

Program goodput also connects directly to the efficiency expressions derived earlier in this book. The compute-time fraction of equation 1.4 divides each device's share of compute time by the distributed step time, and what it measures depends on how that compute time is counted. Counted as the ideal time at peak hardware throughput, the fraction equals the program goodput of a single step. Counted as the measured execution time on a single device, the compute time already contains that device's own losses to memory stalls and kernel dispatch, so the fraction measures only what

spreading the work across devices costs, which is the scaling efficiency. Multiplying it by the single device's MFU then recovers program goodput. The effective fleet throughput in section B.11.2 makes the same split explicit by multiplying peak cluster capacity by single-device MFU, multi-node scaling efficiency, and runtime goodput. The product of those first two terms equals the job's fleet MFU, which is its program goodput, while scheduling goodput is absent only because that expression starts from accelerators the job already holds.

9.2.3 Accounting boundaries

Without sharp boundaries between the three factors, two engineering teams analyzing the same lost accelerator-hour—such as devices held idle by a gang waiting for its remaining ranks, or a step stretched by a straggling peer—will either assign the loss to different factors or subtract it twice. Every hour of fleet capacity must cross three sequential boundaries before it becomes ideal-rate useful work, and a loss is charged to the single factor whose boundary it fails to cross. Table 9.1 pairs each factor boundary with the losses charged there, the telemetry that measures it, the corresponding utilization tier from section 8.12.1, and the chapter that owns the engineering levers to recover it. The final three rows apply the same boundaries to a serving replica.

Table 9.1: Goodput Accounting Boundaries: Each factor ends at one boundary, every loss is charged to the factor whose boundary it fails to cross, and the evidence column names the cheapest record that measures it. Kernel-active time is evidence that helps split runtime and program losses, not a factor. Serving rows apply the same boundaries to a replica, with headroom charged to program goodput and sized by the serving system.

Factor	What ends at its boundary	Losses charged	Evidence	Utilization tier	Lever owner
Scheduling goodput	Fleet capacity-time becomes time held by complete allocations	Unallocated capacity (absent demand, quota limits); devices held by a gang waiting for members; fragmentation; capacity left idle while queued gangs cannot be placed (a job's own queue wait costs only time to result when other jobs hold the capacity)	Scheduler records: allocation intervals with gang-complete timestamps; held capacity per pool; fragmentation index	Allocated, counting complete allocations only	chapter 8
Runtime goodput	Held time becomes time on steps a checkpoint retains	Startup and compilation before the first step; input stalls; synchronous checkpoint pauses; failure detection and recovery; rework after a failure; work lost to preemption; idle reservations	Job timeline: step log with retained-checkpoint markers; restart, preemption, and input-wait logs	Productive, normalized to the job's own fault-free step rate	chapter 7; chapter 4 for input and checkpoint writes
Program goodput	Retained step time becomes ideal-rate time	Exposed communication and synchronization; stragglers; pipeline bubbles; memory stalls; host dispatch; recomputation; contention from asynchronous checkpoint writes	Step traces and device counters; MFU, or MBU for bandwidth-bound decode	Beyond the three tiers	This chapter; chapter 5 and chapter 6 for collectives and layout
Kernel-active time	None; evidence only	None	Device activity counters	Active	None
Scheduling goodput, serving replica	Capacity becomes whole-node replicas held	Devices that cannot form a complete replica; replicas waiting for placement	Scheduler records per replica	Allocated, complete replicas only	chapter 8
Runtime goodput, serving replica	Held replica time becomes ready time	Weight loading, warmup, and compilation; restart after failure; draining before an update	Replica lifecycle log	Productive, at the replica's ideal rate	chapter 7; chapter 10 for cold-start latency

Factor	What ends at its boundary	Losses charged	Evidence	Utilization tier	Lever owner
Program goodput, serving replica	Ready time becomes ideal-rate serving time	Decode below its named ideal; headroom held for bursts; cache fragmentation that limits admission	Per-step traces; tokens served against the named ideal; MBU for decode	Beyond the three tiers	This chapter; chapter 10 decides how much headroom a latency target requires

These boundary definitions refine two of the orchestrator's utilization tiers from section 8.12.1 so that the three factors multiply without overlap. First, scheduling goodput counts only complete allocations. Accelerators held by a gang still waiting for its remaining ranks appear as allocated on a basic cluster dashboard, yet they are charged here as a scheduling loss because no synchronous step can execute on them. Second, the orchestrator's productive tier measures retained work against a declared reference rate. When that reference is the job's own fault-free step rate (with input ready and no interruptions), dividing productive capacity by complete-allocation capacity yields runtime goodput directly. When the reference is hardware peak throughput instead, the same ratio collapses runtime and program goodput into a single number and obscures which boundary lost the capacity. In the orchestration example of 500 GPUs allocated, 350 active, and 280 productive, the productive-to-allocated ratio of 56 percent represents runtime goodput only when the productive tier is normalized to the job's own steady-state step rate.

By contrast, the orchestrator's active tier does not correspond to a goodput boundary. An accelerator whose execution engines are active may be performing required model arithmetic, spinning in a collective while waiting on a straggler, recomputing discarded activations, or re-executing steps lost after a node failure. Those four activities represent useful work, two program goodput losses, and a runtime goodput loss; device activity counters therefore serve as diagnostic evidence for separating runtime from program losses rather than as a goodput factor in their own right.

The boundary between runtime goodput and program goodput also governs how delays inside a step are classified. The duration of a retained step includes any time spent waiting on a slow peer, a pipeline bubble, or a collective operation, but excludes stalls waiting for the storage pipeline to deliver the next batch. Input stalls are carved out and charged to runtime goodput because the data pipeline prevented the step from starting, and the data stall ratio of equation 4.4 in section 4.4 quantifies their share. Every other slowdown within a retained step is charged to program goodput, even when its root cause originates outside the user program.

In short, the accounting boundary determines *where* a loss is counted, whereas the physical cause determines *which system layer* fixes it. Stragglers illustrate why the two must remain distinct. Section 7.9.5 treats a degraded accelerator as a reliability problem and evaluates when replacing the slow node justifies paying for a restart. In goodput accounting, however, a straggler is charged to program goodput: the job holds a complete allocation and retains every step in its checkpoints, yet those steps execute more slowly than the hardware ideal. Recording the straggler in program goodput captures the slowdown where it manifests. Program goodput drops while the straggler is present and recovers once the node is replaced. The remediation decision stays with the reliability controller, which weighs the recovered program goodput against the runtime goodput lost to draining and restarting the job.

Preemption separates cause from charge in the same way. The cluster scheduler triggers the eviction, yet the discarded steps since the last checkpoint and the subsequent reload are charged to runtime goodput because the job held a complete allocation while those hours were spent. By the same rule, a synchronous checkpoint that halts step execution is a runtime loss, whereas an asynchronous checkpoint that competes with the running step for memory or network bandwidth slows the retained step and is charged to program goodput. Ahead-of-time compilation before the first step reduces runtime goodput, while the execution speed of the compiled kernels governs program goodput. Across all of these cases, the three boundaries keep MPG internally consistent, though MPG itself is only one of several metrics in this book that carry the name "goodput."

9.2.4 One word, several measures

Quoting "goodput" without qualification leaves a reader unable to tell whether a job preserved its progress across restarts, used its share of the cluster effectively, met its serving latency targets, or extracted statistical progress from each training example. Across this book, the term appears in four distinct forms, each answering a different engineering question:

- *Run-level goodput* (section 1.5.1.3) measures a single job's rate of useful training progress after discounting time lost to failures and replayed steps.
- *ML productivity goodput (MPG)* measures the fraction of total fleet capacity-time that becomes useful work at the hardware's ideal rate.
- *Service-level objective (SLO) goodput* (section 10.3) measures the rate of completed inference requests that satisfy both time-to-first-token and per-output-token latency targets, which bound the wait before a request's first generated token and the interval between later tokens.
- *Statistical goodput*, which the Pollux scheduler maximizes in section 8.9.4, weights raw training throughput by the convergence progress each training example contributes at the current batch size.

These four measures complement rather than replace one another. SLO goodput applies on top of MPG: a serving fleet that operates at a high fraction of its ideal hardware rate on average can still violate latency targets in the tail if batches are formed too aggressively. Similarly, statistical goodput evaluates progress per example rather than operations per second, so enlarging the global batch beyond the critical batch size can raise program goodput while reducing statistical goodput. Any quantitative efficiency claim must therefore state which goodput measure it reports, which denominator defines its boundary, and which time window it spans. The measurement contract of section 9.3 pins down the workload, boundary, and measurement window so that two goodput numbers can be compared fairly.

9.3 Measuring Goodput

↪ MPG applied across a whole portfolio of jobs, in fleet dashboards and chargeback, is developed in section 12.8.5.

Two measurements of the running Llama 2 70B job taken a week apart can report noticeably different program goodput because they traced different ranks, spanned a framework upgrade, started timing before compilation finished, or averaged across a failure recovery. The job's real efficiency on the cluster need not have changed at all. A goodput factor can guide engineering decisions only when repeated measurements are directly comparable and when a drop can be traced cleanly to the boundary that lost the capacity. Measuring goodput at fleet scale therefore requires three elements: a measurement contract that holds confounding variables fixed, a diagnostic hierarchy that identifies the losing factor before probing its internal mechanism, and a telemetry sampling strategy whose collection overhead the production fleet can afford.

9.3.1 The measurement contract

An apparent speedup between two trials can easily stem from altered input distributions, a warmer machine state, host timers that miss asynchronous kernel execution, or a relaxed convergence target rather than from the optimization under test. Distributed fleets introduce further confounding variables that single-machine benchmarks never encounter: two runs may receive placements with different switch hops, share fabric links with different neighbors, execute under different preemption policies, or span a window too brief to capture routine hardware interruptions. Industry benchmark suites prevent these confounds by locking down the workload, quality target, and timing rules before any system comparison begins (Mattson et al. 2020; Reddi et al. 2019), and production fleet engineering requires the same discipline for every goodput factor it tracks.

Definition 9.3: Measurement contract

Measurement contract is the set of conditions a performance measurement fixes and reports, the workload and its quality gate, the system state and software version, the timing boundary, and the goodput factor measured over a named window with its statistical spread, so that two measurements under the same contract differ only in the change being evaluated.

1. **Significance**: A contract makes the difference between two results attributable to the change under test, so a gain can be priced in accelerator-hours and repeated on another job.
2. **Distinction**: Unlike a benchmark score, which reports one configuration under one set of rules, a contract governs every measurement a team takes, including the before and after of each change. Unlike a monitoring threshold, it states what a measurement must hold fixed to support a claim.
3. **Common pitfall**: A frequent mistake is to fix the workload and hardware while letting the window drift. A runtime goodput measured over a day with no failures describes a better job than the one that runs for weeks on a shared fleet.

A complete measurement contract specifies the five categories in table 9.2, pairing each required record with the misattribution it prevents. The workload envelope locks down the model revision, input and sequence-length distributions, global batch size, and the convergence or output-quality gate that defines valid work. The system-state record captures accelerator count and placement topology, clock and power limits, software, driver, numerical precision, and compiler versions, and—when evaluating scheduling goodput—the cluster pool and scheduling policy in effect. The timing boundary defines which initialization steps fall inside the timed region and how the host synchronizes with asynchronous device streams. Finally, the statistical report identifies the goodput factor, its denominator, and its calendar window alongside the trial count, central tendency, dispersion, step-time percentiles, and quality verification. Percentiles belong in the report because a mean step time can hide the slow tail episodes of an intermittent straggler. The fifth category, load generation, applies to serving measurements and must keep offering requests at the specified rate while the replica stalls, because a load generator that waits for each response eases off during a stall and hides it (chapter 10).

Table 9.2: Reproducible Performance Measurement Contract: Every goodput result fixes workload, correctness, system state, timing, and statistical reporting, and names the factor and window it measures, before two implementations or configurations are compared.

Contract field	Evidence to report	Failure prevented
Workload and correctness	Model revision; representative input and sequence-length distributions; global batch or concurrency; quality or output-equivalence gate	Comparing different work or accepting an invalid numerical shortcut
System state	Device count and topology; clocks and power limit; software, compiler, driver, precision, and kernel versions; for scheduling goodput, the pool and scheduler policy	Attributing a hardware, software, or policy change to the optimization
Timing boundary	Included setup; device synchronization; warmup; cold-start and steady-state definitions	Under-timing asynchronous execution, or hiding startup and compilation inside a steady-state number
Statistical report	Factor measured and its denominator; window; trial or window count; median or mean; dispersion; step-time percentiles	Promoting a favorable single run, or multiplying factors measured over different windows or overlapping denominators
Load generation (serving)	As specified for serving measurements in chapter 10	Load that eases off while the replica stalls

The required measurement window varies sharply across the three factors because each boundary loses capacity on a different timescale. Program goodput captures inefficiencies that repeat on every training step or decode iteration, so a window of several hundred warmed steady-state steps estimates it reliably. Runtime goodput captures episodic events such as node failures and rollbacks to the last checkpoint. A window shorter than the mean time between cluster interruptions can miss those events entirely and report a misleadingly high factor, so its window must span days to weeks at fleet scale (chapter 7). Scheduling goodput tracks cluster contention, which fills queues during peak working hours and drains them overnight or between project deadlines, requiring a window that spans full demand cycles. Multiplying the three factors yields a job's true MPG only

when all three cover the same calendar window, even when program goodput within that window is estimated by sampling representative retained steps.

The timing boundary likewise determines whether startup overhead is attributed to runtime goodput or program goodput. End-to-end trials that start the clock before graph compilation and cache warmup measure the delay before the first productive step, which belongs to runtime goodput, whereas warmed steady-state trials isolate the per-step execution rate, which belongs to program goodput. Mixing cold-start and warmed runs can credit a code change for a speedup when the second trial simply hit a warm compiler cache; the measurement contract therefore reports cold-start duration and steady-state step time separately. When a more aggressive compiler pass speeds steady-state steps at the cost of a longer initial compile, it trades runtime goodput against program goodput—a trade-off that section 9.9 prices explicitly.

Industry benchmark suites illustrate how a strict measurement contract operates in practice. MLPerf Training measures wall-clock time to train a reference model to a specified validation target, and MLPerf Inference measures throughput under explicit latency bounds across standardized serving scenarios, evaluating complete hardware-software systems rather than isolated chips (Mattson et al. 2020; Reddi et al. 2019). Closed-division submissions hold the model architecture and optimizer fixed so that systems can be compared directly, while open-division submissions relax those constraints to showcase algorithmic and systems co-design. Because a benchmark submission runs on a dedicated, pre-warmed cluster allocation without multi-tenant queueing or long-run preemption churn, it primarily measures program goodput and reveals very little about scheduling or runtime goodput (perspective 9.1).

Recording exact software versions in the system-state contract is essential because software upgrades alone can shift delivered throughput by a large multiple on identical silicon. Between September 2023 and April 2025, NVIDIA submitted the same eight-GPU DGX H100 node to four consecutive MLPerf Inference rounds while upgrading its inference stack from TensorRT 9.0 to 10.8 (Reddi et al. 2019; MLCommons 2025a). Across those four rounds, offline throughput rose 2.7× on GPT-J, 1.4× on Llama 2 70B, and 1.2× on Stable Diffusion XL (figure 9.3), whereas ResNet-50 and BERT—whose kernels were already mature—remained flat. Any before-and-after experiment that straddles a framework or compiler upgrade therefore confounds the upgrade with the optimization under test.

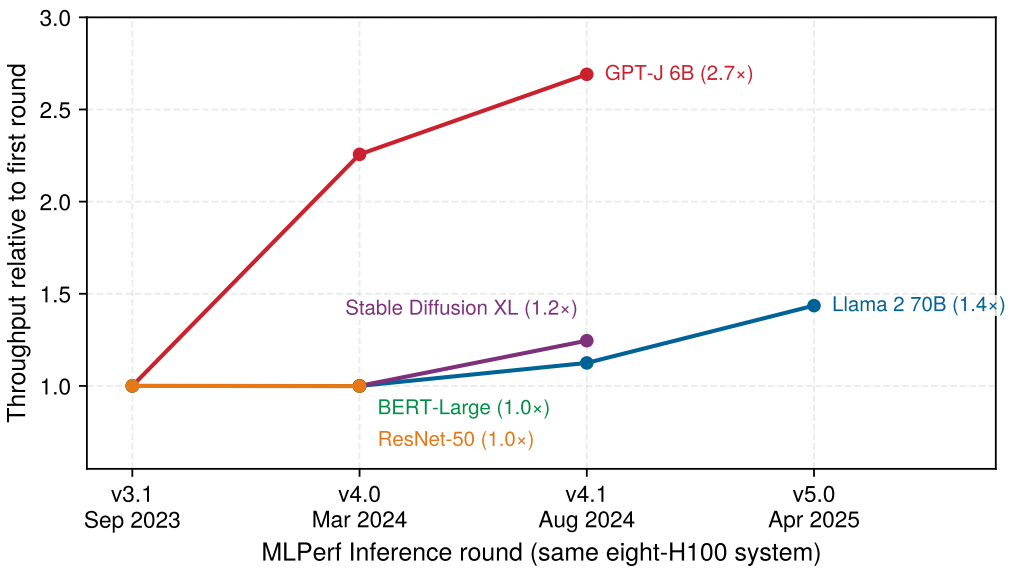

Figure 9.3: Software-Only Gains on Fixed Hardware: Offline throughput of one NVIDIA DGX H100 system (eight H100 SXM GPUs with 80 GB of memory each) across four MLPerf Inference rounds, relative to each benchmark's first submission on that system. The hardware is identical in every round; the software stack changed from TensorRT 9.0 to 10.8. Llama 2 70B uses the 99.9 percent accuracy target. Data from MLCommons (2025a).

9.3.2 Localizing the lost factor

Even when a contract-controlled measurement confirms a genuine slowdown, knowing that a training run fell behind its schedule does not reveal which goodput boundary lost the capacity. An engineering team that opens a kernel profiler first will always find some exposed gap inside a step. The team can easily spend a week shaving milliseconds off a collective while most of the missing capacity was lost waiting for a complete allocation or recovering from unlogged restarts. Effective localization instead inspects telemetry from the cheapest, broadest record down to the most expensive, narrowest trace, matching each record to one boundary of table 9.1.

At the top of that hierarchy sit scheduler records. For every submitted job, the cluster scheduler logs submission time, individual device assignments, the timestamp at which the minimum gang became complete, final release, and the resource pool, quota, and priority class under which the job ran. Pool-level logs add unallocated capacity and the fragmentation index from section 8.12.1. Because the control plane already writes these records for admission control and chargeback, they carry zero extra instrumentation overhead and cover every job in the cluster. Scheduler records establish scheduling goodput directly and supply the all-allocated held time that forms the denominator of runtime goodput. What they cannot see is anything that happens *inside* a complete allocation: a wedged job that holds a full gang without executing a single step looks identical in scheduler logs to a job running at peak efficiency.

The second rung is the job timeline, the structured lifecycle log emitted by the job itself between its gang-complete timestamp and its release. The timeline records when the first training step began, the completion timestamp of every step, which steps each committed checkpoint retained, every failure, restart, and preemption event, and the time each step spent waiting on the input pipeline. Because it stores only a few scalar fields per step or lifecycle event, the timeline is compact enough to retain for every production job. It settles runtime goodput. The interval before the first step measures startup and compilation, pauses at checkpoint commits measure synchronous write overhead, steps repeated after a rollback measure rework, and the input-wait counter yields the data stall ratio of equation 4.4. Because the timeline also logs the wall-clock duration of every retained step, it flags a program goodput loss whenever retained steps run longer than their baseline target. It cannot distinguish, though, whether that slowdown came from exposed communication, memory stalls, or a straggling rank.

The third rung consists of step traces, the only telemetry that looks inside an individual step. Over a selected time window and set of ranks, a step trace captures every kernel launch, collective operation, memory copy, and host dispatch interval alongside hardware performance counters such as achieved HBM bandwidth. That detail partitions a retained step into computation, exposed communication, peer or pipeline waits, memory stalls, and host dispatch overhead, pinpointing which term of the Fleet Law (local work, communication, or synchronization; section 9.6.1) an optimization must shorten. Because full tracing generates gigabytes of data and perturbs execution timing, it is reserved for short windows on sampled ranks; a clean step trace taken during a healthy hour says nothing about whether the job lost half its week to restarts.

Table 9.3: Records That Localize a Goodput Loss: Each record matches one accounting boundary, and the cheaper records come first because they cover the whole fleet continuously. A loss is traced to its factor before its mechanism is studied, and the step trace is collected only for a program loss the job timeline has already detected.

Record	Coverage and cost	Establishes	Cannot establish	Factor it settles
Scheduler records	Every job and pool, continuously; written for admission and chargeback	Capacity not held by complete allocations; queue wait; gang-complete times; held time per job	Anything inside a held allocation, including whether any step ran	Scheduling goodput; the denominator of runtime goodput
Job timeline	Every job, for its whole life; a few fields per step or event	Startup, input waits, checkpoint pauses, restarts, preemptions, rework; retained step time; step-time distribution	Why a retained step takes as long as it does	Runtime goodput; detects but does not explain program loss

Record	Coverage and cost	Establishes	Cannot establish	Factor it settles
Step traces	Selected ranks and short windows; large volume and some perturbation	Intervals inside a step (computation, exposed communication, peer and stage waits, memory stalls, host dispatch)	Losses outside the traced window (scheduling waits, restarts, checkpoint pauses, intermittent episodes)	Program goodput, and the Fleet Law term to shorten

Reading these three records from top to bottom (table 9.3) is a methodological requirement, not merely a cost-saving habit. Each rung settles one goodput factor and hands the remaining unexplained gap to the rung below, so engineers collect step traces only after scheduler logs and the job timeline confirm that retained steps themselves are running slowly. Skipping rungs leads to two classic diagnostic blunders. Starting with a step trace can pronounce a job healthy while it loses a third of its held hours to unexamined restarts. Starting and stopping at queue logs can blame the cluster scheduler for a week-long slowdown that the job timeline would immediately trace to a compiler regression. In both failure modes the local data are accurate, yet the diagnosis fails because the chosen telemetry cannot observe the boundary where the loss occurred.

For these three records to enforce the accounting boundaries of section 9.2.3, their instrumentation must align with those boundaries. If a job timeline logs only total step duration without separating data-loader stalls, every input wait is folded into step time and a runtime goodput loss masquerades as a program goodput loss. Similarly, because a single straggling accelerator stretches synchronous step time across every peer in its collective group, identifying which rank caused the slowdown requires lightweight per-rank step timers inside the job timeline.

Our running Llama 2 70B training job illustrates how the three rungs work together. Suppose that after one week on the shared cluster, the job has retained substantially fewer steps than predicted by a plan that assumed full-week allocation at the single-node step rate. Consulting the scheduler records first reveals that the job held partial node allocations before its full 16-node gang assembled—a scheduling loss quantified in section 9.4.3. Reading the job timeline from that gang-complete timestamp next isolates the initial startup delay, synchronous checkpoint pauses, steps re-executed after failures, and a modest input-stall fraction, which section 9.5.2 composes into runtime goodput. The same timeline also shows that every retained step ran longer than the single-node baseline from day one, proving that part of the shortfall lies in program goodput. Only at that point does the team trigger step traces on one data-parallel group over a brief window. Because per-rank timeline counters show no single outlier rank, the team inspects intervals common to all ranks and finds cross-node synchronization and collective waits with no computation beside them on every step, the exposed communication analyzed in section 9.6.2. Each rung answers the exact question that the others cannot (figure 9.4).

Diagnosing the eight-GPU serving replica follows the exact same three-rung progression. Cluster scheduler logs measure how long the replica waited for a full unfragmented node, the replica lifecycle log tracks weight loading, warmup, ready service, and restarts to establish runtime goodput, and decode step traces—evaluated against peak HBM bandwidth via MBU—measure how closely ready execution approaches the hardware ideal.

9.3.3 Fleet profiling by sampling

The first two rungs of the localization ladder scale effortlessly with cluster size because a scheduler log entry or a step-timeline record occupies only a few bytes whether a job spans one node or a thousand. Step traces do not scale in that way. When a job expands from 8 GPUs to 1,024 GPUs, multi-node bottlenecks emerge that no single-node trace can capture—such as fabric-crossing collectives and ranks stalled behind a distant straggler—so profiling must span many ranks at once. Yet if a 5-second Nsight Systems trace on one GPU produces 500 MB, capturing that same 5-second window across 1,024 GPUs generates roughly 512 GB of telemetry. No production cluster can ingest or store traces at that rate continuously, and logging every kernel launch perturbs the very step timing under study (perspective 9.2). Deciding *which* ranks and time windows to trace is therefore a core part of fleet measurement.

Production fleets resolve this tension through hierarchical profiling, splitting telemetry between lightweight, always-on counters that detect an efficiency drop and targeted traces that diagnose

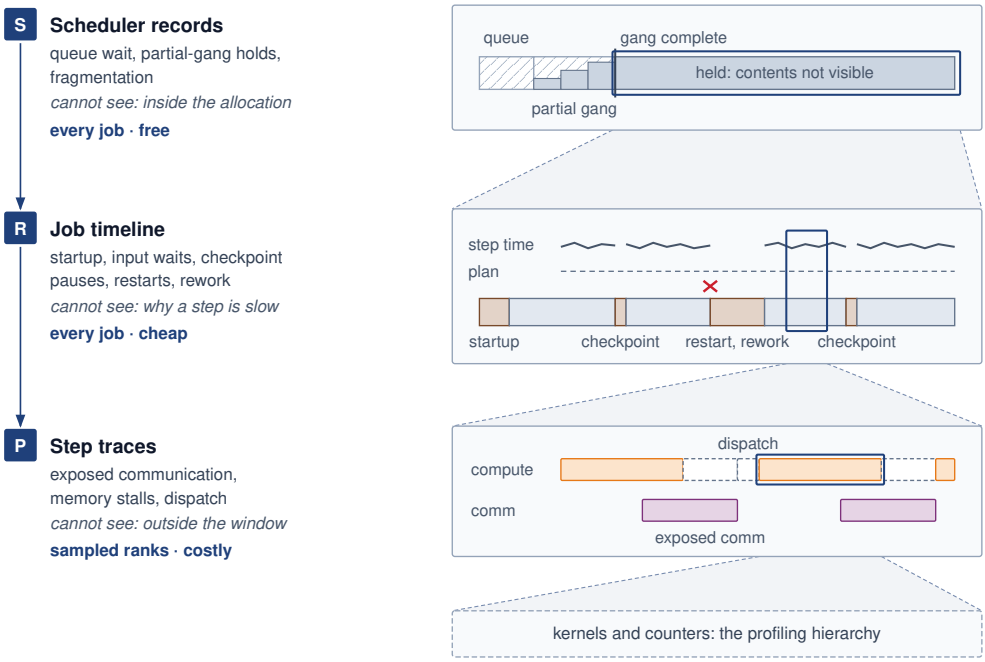

Figure 9.4: Localizing a Lost Factor: The three records are read from the top, cheapest and most complete first, and each settles one factor of goodput. Scheduler records cover every job and settle scheduling goodput, but they see a held allocation only as a closed block. The job timeline settles runtime goodput and shows that retained steps run slower than planned without showing why. Step traces, collected for sampled ranks over a short window, settle program goodput, and their highlighted window expands into the profiling hierarchy of figure 9.6. Widths are schematic, not measured.

its cause. Within a strict overhead budget, the always-on layer continuously records step duration, retained-step MFU, per-rank step timers, and hardware health telemetry (HBM bandwidth, power draw, and die temperature) across every active job. Once those counters flag that a goodput factor has degraded and narrow down where, targeted profiling captures a short trace window from specific ranks chosen to test a hypothesis. Comparing a suspected straggler directly against a healthy peer in the same collective group isolates the stall far more cleanly than tracing an arbitrary rank. ByteDance's MegaScale system[4] exemplifies this workflow. Lightweight CUDA-event timers record critical operations on every rank throughout production training, summarize per-rank latencies as a cluster-wide heat map, and drill into a synchronized timeline only for the parallel group that exhibits an anomaly (Z. Jiang et al. 2024). Below the step-trace rung, diagnosis then descends through framework operators, kernel timelines, and hardware performance counters along the profiling hierarchy of figure 9.6 (section 9.6.3).

To inspect inside steps continuously without paying for full traces, fleets rely on statistical profiling. A sampling policy selects a small fraction of kernels, steps, or ranks for detailed timing, and aggregating those samples reconstructs where step time is spent across the fleet. Google-Wide Profiling pioneered this approach across warehouse-scale computers by continuously sampling a small fraction of machines and aggregating the results into fleet-wide execution profiles (Ren et al. 2010). Applied to an ML accelerator fleet, fleet-wide aggregation ranks kernels and collectives by their share of total retained step time across all jobs. A collective primitive that consumes a modest slice of every job's step often represents a far larger fleet-wide optimization target than a specialized kernel that dominates a single workload.

Statistical sampling is only as trustworthy as its stratification. A fleet sampling policy must stratify across job sizes and model architectures, pipeline stages and parallel groups, warmup and steady-state phases, and hours of the day. Sampling only rank 0 (which often performs extra logging and checkpoint coordination) or only the opening minutes of a job characterizes a special rank or

4 **MegaScale stragglers**: The heat map of per-rank forward and backward latency showed about 0.5 percent of machines running substantially slower than the rest, and excluding them made peak MFU consistent across runs (Z. Jiang et al. 2024). Separately, a few hosts took about 10 percent longer on the same forward computation, and removing them raised MFU by about 0.7 percent. A fraction of a percent of machines thus set the step time of every rank that synchronized with them, a program loss that only per-rank timing on every step could locate.

phase rather than the fleet. Furthermore, because rare anomalies that occur once every several thousand steps can easily slip between periodic random samples, continuous sampling is paired with threshold triggers that automatically dump a detailed trace whenever an always-on timer crosses an anomaly bound.

Intermittent stragglers illustrate why triggered capture is essential. An accelerator can slow down for brief episodes due to thermal throttling, recoverable memory errors, transient fabric congestion, or host OS pauses such as page reclamation. Because a synchronous step completes only when its slowest participant finishes, a single rank running slower during an episode delays its entire collective group by up to that same fraction. The job's cumulative program goodput loss scales with both the slowdown severity and the episode frequency. A periodic trace captured between episodes sees completely healthy execution. Always-on per-rank step timers catch the episodes as latency spikes, and crossing a duration threshold immediately triggers synchronized traces on both the lagging rank and a healthy peer. Correlating those triggered traces with device telemetry, switch counters, and host kernel logs then separates the rank that originated the stall from the peers that merely waited at the collective barrier.

9.3.4 Detecting regressions

↪ How fleet telemetry is organized into monitoring levels that route an alert to the team that owns it is developed in section 12.6.3.

A production fleet changes continuously even when a model team leaves its training script untouched. Framework and compiler releases alter the kernels a model dispatches, driver and firmware updates modify clock management and collective algorithms, storage migrations shift input and checkpoint bandwidth, and scheduler policy updates reshape how allocations form. Any of those changes can erode a goodput factor across every job it touches—often only at multi-node scale, where a slight increase in peak activation memory triggers allocator defragmentation on one rank or a longer host pause stalls an entire synchronous gang. Systematic regression detection applies the measurement contract to every infrastructure or software change and monitors each goodput factor in the telemetry record that settles it (table 9.4).

Table 9.4: Regression Test Tiers by Goodput Factor: Program regressions are caught by canaries and fleet baselines of MFU and step time; runtime and scheduling regressions are caught only by baselines kept over long windows in the job timelines and scheduler records, because their losses are rare per step and appear across many jobs.

Testing tier	Factor	What it measures	Regression caught
Small-scale canaries	Program	Step time and MFU on one node and on a small multi-node allocation, establishing a scaling efficiency curve	Code changes that look harmless on one GPU but bend the curve before full-fleet launch
Fleet baseline comparison	Program	Every production run's MFU of retained steps against the reference baseline for that architecture and software version	Architecture, compiler, or configuration changes that reduce useful fleet work
Gray failure detection	Program	Distribution of step times and per-rank step timers across the fleet	Straggler nodes; one node slowed by some fraction can lower synchronous data-parallel MFU by about that fraction
Startup and recovery baselines	Runtime	Time from gang completion to first step; checkpoint pause per commit; restarts, rework, and data stall ratio per day	Compiler or framework upgrades that lengthen compilation, storage changes that slow checkpoint writes or input, recovery automation that stops working
Allocation baselines	Scheduling	Queue wait to gang completion; capacity held in incomplete gangs; fragmentation index per pool over whole demand cycles	Scheduler policy or job-template changes that strand capacity or delay complete allocations

The program-goodput tiers catch most regressions introduced by model or kernel changes within minutes, because any per-step slowdown repeats on every iteration and manifests on a small multi-node canary. The runtime and scheduling tiers catch regressions that short canary runs cannot observe: a compiler release that adds minutes to graph compilation leaves steady-state step time untouched, and a quota policy that strands nodes in a shared pool surfaces only across a full diurnal demand cycle. Maintaining per-software-version and per-pool baselines under the same measurement contract allows engineers to tie any shift in scheduling, runtime, or program goodput

directly to the release or policy change that introduced it. Benchmark runs miss the same runtime and scheduling losses that short canaries miss, so a cluster sized from benchmark numbers inherits every loss the benchmark omitted.

↩ Gray failures, the silent degradations that surface as stragglers rather than crashes, are established in section 7.9.5.

Systems Perspective 9.1: What a benchmark run measures

Benchmark results, including MLPerf submissions, describe one tuned run under a specified protocol on a dedicated allocation. In goodput terms they measure program goodput, and useful work through the quality target, with scheduling goodput equal to one by construction and few of the interruptions a long production run accumulates. They are a reference for the program factor, not a capacity guarantee.

Production adds what the reference leaves out. Startup, checkpoint pauses, failures and rework, preemptions, and input waits are runtime losses, monitoring and contention with neighbors can lower program goodput further, and waiting for complete allocations is a scheduling loss no benchmark sees. As a planning scenario, if useful throughput is 10–20 percent below the rate that implied 30 days of training, the same work takes about 33.3–37.5 days. Section 9.5 sizes the runtime share of that gap for a well-run job.

Detecting a regression and verifying an optimization are the same measurement performed in opposite directions. Every candidate change must be evaluated at the accounting boundary it targets—in retained step time and MFU for a program change, in the job timeline for a runtime change, and in scheduler logs for an allocation policy change—rather than in an isolated microbenchmark. Doubling the speed of a single kernel saves at most half of that kernel's share of the step, and saves zero wall-clock time if the kernel already ran hidden behind an exposed collective. Because an optimization at one layer can degrade another factor, all three goodput factors are re-measured under the contract before promotion. Even a change that appears purely local can exhaust a shared hardware resource when deployed across thousands of nodes simultaneously, which is why production fleets gate infrastructure changes behind staged rollouts (section 12.5.5). One of the clearest postmortems of a local cost turning into a fleet-wide outage comes from edge networking rather than ML training.

War Story 9.1: The regex that saturated the edge (2019)

Context: Cloudflare's web application firewall evaluates rules on HTTP/HTTPS traffic across its global edge (Cloudflare 2019).

Mechanism: A rule deployed on July 2, 2019 caused excessive regular-expression backtracking and drove the CPUs serving HTTP/HTTPS traffic to nearly 100 percent utilization across the network. A previously removed CPU-use protection, missing performance tests, and a global rollout allowed the local computational cost to become a fleet-wide failure.

Impact: The resulting 27-minute outage disrupted proxying, content delivery, and firewall service. Traffic through the network dropped by about 82 percent, and users received 502 errors.

Fix: Disabling the firewall globally restored traffic; the offending change was then removed and the firewall re-enabled after testing. The postmortem reported restoring CPU-use protection and plans for performance tests, staged rule deployment, and a regular-expression engine with runtime guarantees.

Systems lesson: Correct outputs do not establish acceptable execution cost. A change that leaves every result correct can still lower goodput across the whole fleet at once. Performance-sensitive changes to an ML fleet, such as a compiler upgrade, a kernel library, or a new input shape, are therefore measured under the contract, staged on a small share of the fleet, and promoted only when the factors they touch hold.

A measurement contract that locks down each factor's window, a three-rung telemetry ladder that localizes losses from the cheapest record downward, and continuous baselines that catch regressions

together make each goodput boundary observable in production. The first boundary a workload encounters is the formation of a complete allocation, where fleet capacity can be lost before a single step begins.

9.4 Scheduling Goodput

Partially assembled gangs, scattered node fragments, and draining reservations waste accelerator hours *before* any workload holds a complete allocation, whereas zombie jobs, input starvation, exposed collectives, and fragmented key-value caches waste capacity inside an allocation a job already holds. A cluster orchestrator's general health dashboards surface both classes of waste together, but scheduling goodput counts only the pre-allocation losses; fixing the post-allocation losses leaves scheduling goodput unchanged. Separating pre-allocation losses from post-allocation losses, and quantifying the pre-allocation share, is what scheduling goodput adds to the orchestrator's own view. The policy levers that recover those scheduling losses (atomic gang admission, backfill, bin packing, topology-aware placement, quota lending, and preemption policy) are developed in chapter 8.

9.4.1 What scheduling goodput counts

A rigid distributed training job holding half of its requested accelerators while waiting for the rest completes zero steps, not half as many, yet a naive utilization counter credits every assigned device as busy. Scheduling goodput eliminates that distortion by taking the ratio of two capacity-time integrals over the same calendar window. The denominator is the in-service capacity-time offered by the pool: its allocatable accelerators multiplied by the window duration. The numerator is the all-allocated capacity-time held by complete allocations, summed across jobs as each job's device count multiplied by the interval from the moment its minimum gang became complete until release. A complete allocation corresponds to the minimum viable execution group of section 8.4.3—every accelerator, interconnect domain, and co-scheduled host resource required to execute one synchronous step. For a rigid job, that group is the full requested gang; for an elastic job, it is the active gang size between two reconfiguration boundaries, and any accelerators the job relinquishes re-enter the numerator only when another complete allocation claims them. Because nothing inside the allocation enters this ratio, a wedged job that holds a complete gang without running a step achieves the same scheduling goodput as a job executing at peak FLOP/s.

Every device-hour of in-service capacity outside a complete allocation is a scheduling loss, regardless of its operational cause. Capacity sits unassigned when offered demand falls below pool size, when waiting jobs have exhausted their team quotas, or when free accelerators are scattered across racks in fragments too small for any queued gang. Capacity is assigned yet incomplete when a gang holds an initial subset of nodes while waiting for the remainder, when two partial gangs deadlock over the last free nodes, when freed accelerators sit idle in a reservation draining for the head-of-line gang, or during brief container teardown and setup transitions between tenants. Because an incomplete gang cannot advance synchronous training, section 9.2.3 restricted the orchestrator's allocated tier to complete allocations alone.

Accelerators removed from service for hardware repair or scheduled maintenance are excluded from the scheduling goodput denominator because no scheduling policy can place work on broken nodes. Those out-of-service hours are tracked alongside MPG as cluster availability, keeping hardware downtime visible without blaming the scheduler. If a report instead divides by total installed silicon, hardware unreliability is folded into scheduling goodput. The measurement contract (section 9.3.1) therefore requires stating whether the denominator is in-service or installed capacity so that two fleet reports are never compared on mismatched baselines.

When scheduling goodput is low, scheduler logs distinguish between two fundamentally different causes that call for opposite remedies. If accelerators sat unallocated while the job queue was empty, the pool exceeds current demand—a capacity-planning and quota-lending problem. If accelerators sat unallocated or trapped in partial gangs while runnable jobs waited in the queue, placement constraints, rigid quotas, or fragmentation stranded the hardware—a scheduling-policy problem. Recording queue depth alongside every interval of unheld capacity separates the two immediately.

Because unallocated fragments, partial gangs, and draining reservations belong to the shared pool rather than to any single workload, scheduling goodput is a property of the resource pool. Our running Llama 2 70B job inherits the scheduling goodput of the pool from which it draws its nodes.

↪ Sizing a fleet against forecast demand, the remedy for capacity that no queued work wants, is developed in section 12.9.5.

A job's own time spent waiting in the queue is a separate metric. Queue wait reduces scheduling goodput only during intervals when pool accelerators actually sat outside complete allocations. While other complete jobs occupy the cluster, a queued job's wait costs zero fleet goodput and instead lengthens that job's time to result, a trade-off priced in section 9.9. As the queueing analysis in section 8.3 demonstrated, driving pool utilization and scheduling goodput toward one causes queue waiting times to rise steeply.

9.4.2 Attributing scheduler losses

Defining scheduling goodput strictly at the complete-allocation boundary means that not every failure mode managed by the orchestration team is a scheduling loss. If a zombie job were subtracted from scheduling goodput simply because orchestrator lease timeouts clean it up, those wasted hours would be deducted from capacity that the scheduling numerator already counted as a complete allocation, penalizing the same hours twice. Attributing a capacity loss to whichever team owns the remediation lever therefore breaks the multiplicative goodput identity. Whereas the operational debugging workflow in section 8.12 groups utilization anti-patterns by their observable symptoms, goodput accounting classifies each loss pattern by the *first boundary its capacity fails to cross*. Capacity outside a complete allocation reduces scheduling goodput; capacity inside a complete allocation that produces no checkpointed step reduces runtime goodput; and capacity in retained steps that execute below the hardware ideal reduces program goodput. Because each lost interval stops at the first boundary it fails, every device-hour is charged once. Applying that boundary test to the orchestration anti-patterns shows that the system layer responsible for fixing a loss often differs from the goodput factor that records it (table 9.5).

↩ The utilization anti-patterns and the evidence that diagnoses each one are established in section 8.12.3.

Table 9.5: Attributing Orchestration Losses to Goodput Factors: Each utilization loss pattern of fleet orchestration is charged to the first accounting boundary its capacity fails to cross. Scheduling goodput receives only the losses that occur before a complete allocation exists, and a topology mismatch lands in scheduling or program goodput depending on whether the job waited or ran. The factor a loss lowers and the layer that removes it are recorded separately.

Loss pattern	What the capacity is doing	First boundary missed	Charged to	Layer that removes it
Partial-gang hold, including hold-and-wait deadlock	Assigned to a rigid job still waiting for members, so no step can run	Complete allocation	Scheduling goodput	Orchestration (atomic group admission)
Reservation for a waiting gang	Freed devices kept idle so the gang at the head of the queue can assemble	Complete allocation	Scheduling goodput	Orchestration (backfill under enforced runtime bounds)
Fragmentation and capability mismatch	Free devices scattered so that no queued gang fits, or of a type no queued job requests	Complete allocation	Scheduling goodput	Orchestration (bin packing, capability matching)
Topology mismatch, waiting	Held free while a compact placement assembles	Complete allocation	Scheduling goodput	Orchestration (placement policy)
Topology mismatch, running	Running across slower links than a compact placement would use	Ideal rate	Program goodput	Collectives and parallel layout; orchestration (placement policy)
Idle reservation	Reserved block held as an allocation while no job takes steps	Retained step	Runtime goodput	Orchestration (quota reclamation and lending)
Zombie job	Complete allocation kept after execution has stopped	Retained step	Runtime goodput	Fault tolerance (failure propagation); orchestration (leases, progress monitoring)
Preemption	Steps since the last checkpoint discarded; state reloaded and warmed up	Retained step	Runtime goodput	Orchestration (preemption budgets); fault tolerance (checkpoint cadence)
Data starvation	Steps waiting for their next input batch	Retained step	Runtime goodput	Data storage (input pipeline)

Loss pattern	What the capacity is doing	First boundary missed	Charged to	Layer that removes it
Communication bottleneck, including noisy neighbors	Retained steps expose collectives that computation does not hide, or that a co-located job's traffic slows	Ideal rate	Program goodput	Collectives and parallel layout; orchestration (isolation)
Cache fragmentation limiting admission (serving)	A ready replica admits fewer concurrent requests than its node could serve	Ideal rate	Program goodput	Serving (paged cache allocation)

Idle accelerators appear in two different goodput factors, depending on their allocation state. When a partial gang, a stranded fragment, or a draining head-of-line reservation leaves accelerators idle *outside* any complete allocation, the loss belongs to scheduling goodput. By contrast, when a team holds a static capacity reservation (section 8.11.4) as a complete allocation while running no steps on it, that capacity crosses the scheduling boundary and fails at the runtime boundary, registering as a runtime goodput loss. Reclaiming an idle quota block and returning it to the shared pool therefore shifts the accounting before it improves fleet output. Whatever fraction of the reclaimed block a borrowing team puts to work is recovered, while any fraction that finds no borrower sits free in the pool and converts from a runtime loss into a scheduling loss. An orchestration policy can shift waste from one boundary to another without increasing useful work; only the product $\eta_{\text{scheduling}} \times \eta_{\text{runtime}} \times \eta_{\text{program}}$ reveals whether the fleet came out ahead.

In every row of table 9.5, the accounting charges the specific time interval that exhibits the loss pattern rather than tainting a job's entire run. A zombie job wastes its complete allocation from the moment progress halts until the lease expires and releases the nodes (section 8.4.4). An input-starved job's data-loader stalls are charged to runtime goodput from the job timeline's input-wait counter, while the active execution intervals between stalls pass into retained step time and are judged on their program goodput.

Finally, a communication or placement bottleneck on a running job lies entirely past the first two boundaries, so it is charged to program goodput, and a targeted step trace isolates the exposed collective in the Fleet Law's communication term (section 9.6.2). Once those post-allocation rows are routed to runtime and program goodput, the pre-allocation rows that remain can be measured directly from the scheduler's logs.

9.4.3 Evidence, size, and the serving replica

Enforcing the complete-allocation boundary in production hinges on a single timestamp in the scheduler logs, the top rung of section 9.3.2. The cluster scheduler must explicitly record the instant a gang allocation becomes complete, as Kubernetes `PodGroup` admission does in section 8.5.2, rather than inferring completion from the job's first training step. Inferring gang completion from the first step would sweep process launch and compilation into the pre-allocation window, misclassifying a runtime startup loss as a scheduling loss. Likewise, if a scheduler logs only individual node bindings without a gang-complete event, analysts can reconstruct completeness for rigid jobs from their requested gang size, but cannot do so reliably for elastic jobs whose valid size changes over time.

When that gang-complete timestamp is present, scheduler logs provide everything needed to compute scheduling goodput: summing each job's device count over the interval from gang completion to release yields the numerator, summing partially assigned gangs from first device assignment to gang completion quantifies partial-hold waste, and tracking unallocated devices alongside queue state accounts for the remainder. The utilization breakdown from section 8.13 (the pitfall of treating all allocated GPU-hours as productive) illustrates how these boundaries partition a shared cluster pool of the kind that hosts our running Llama 2 70B job.

> Napkin Math 9.2: Scheduling goodput of a shared pool
>
> **Problem**: In a shared fleet's snapshot, 80 percent of in-service capacity is assigned to jobs, 15 percent is held by groups still waiting for their remaining members, and another 5 percent sits

in complete allocations waiting for input. What is the pool's scheduling goodput, and how far can atomic group admission alone raise it?

Variables:

- Assigned share 0.80, which includes the incomplete-group share 0.15.
- Input-wait share 0.05, which lies inside complete allocations and is therefore a runtime loss, sized in section 9.5.

Math:

- Scheduling goodput: $\eta_{\text{scheduling}} = 0.80 - 0.15 = 0.65$.
- Scheduling loss: $1 - 0.65 = 0.35$, the 0.20 no job holds plus the 0.15 held by incomplete groups.
- Ceiling for removing partial holds: group admission returns the 0.15 to the pool as free capacity, so $\eta_{\text{scheduling}}$ rises at most to $0.65 + 0.15 = 0.80$, and only if other jobs' complete allocations take every freed device.

Result: The pool's scheduling goodput is 0.65, so 35 percent of its capacity is lost before any step runs, and removing partial holds alone can raise it no higher than 0.80.

Systems insight: The scheduling loss of 0.35 is more than three times the at most 0.10 of held capacity that Llama 3 405B lost to all its interruptions on a dedicated cluster. Because the factors multiply, every later gain in a job's runtime or program goodput reaches this pool at 0.65 of its value on a dedicated cluster.

This constructed pool snapshot—which also anchors the whole-fleet synthesis in section 9.8—shows why atomic group admission is necessary but not sufficient. Releasing partial-gang holds raises scheduling goodput only insofar as backfill and bin packing refill the freed accelerators with complete gangs, and pushing past the 0.80 ceiling requires runnable demand for the remaining unallocated devices. Queue logs distinguish whether that unheld capacity is stranded by fragmentation and rigid quotas, which scheduling policy can recover, or lacks runnable demand altogether, which only quota lending or capacity planning can address.

A single fleet-wide average also conceals how scheduling policy distributes goodput across job sizes, which is why production fleets track scheduling goodput by size cohort. Across Google's TPU fleet, scheduling goodput exceeds 95 percent for every size cohort, peaking at the smallest and largest jobs while dipping for medium-sized jobs (Wongpanich et al. 2025). Every one of those measured cohorts sits far above the constructed snapshot of this section, which illustrates the scheduling losses rather than describing a measured fleet. Preemption policy explains that U-shaped curve. Evicting a multi-pod training run would discard a long initialization and significant post-checkpoint work, so the scheduler shields the largest gangs from preemption. At the other extreme, single-node jobs finish quickly and pack easily into small gaps when preempted, so the scheduler uses them to defragment the pool and concentrates evictions on medium-sized gangs instead. Raising the scheduling guarantee for one size class therefore lowers it for another—a policy trade-off that a single fleet-wide average hides.

The eight-GPU serving replica crosses the scheduling boundary as a single indivisible gang: one full node across which tensor parallelism partitions the model weights. Because each replica requires a whole node, serving pools are sensitive to the intra-node fragmentation that small training or evaluation jobs leave behind. A pool can have dozens of free GPUs scattered across partial nodes and still have zero whole nodes available. Both a replica waiting in the queue for an unfragmented node and dedicated serving nodes sitting without a placed replica count as scheduling losses. Once placed on a node, however, the replica holds a complete allocation even before it can answer requests: weight loading, warmup, and graph compilation are runtime losses, while spare capacity held on a ready replica for traffic bursts is a program loss. Placement records, replica lifecycle logs, and decode traces separate the three. For a replica as for a training job, scheduling goodput ends where a complete allocation begins, which leaves open how much of that held time becomes retained work.

9.5 Runtime Goodput

A training job that loses held time to input stalls, synchronous checkpoint pauses, hardware failures, and forced preemptions encounters each loss in a different unit. Chapter 4 expresses input stalls as a fraction of step time, chapter 7 expresses checkpointing and recovery overhead against fault-free training time, and chapter 8 prices preemption in wasted minutes per eviction event. Runtime goodput—the fraction of a job's all-allocated held time that becomes retained progress—places all of those losses on a single common denominator. Expressed as shares of held time, every interruption can be compared directly so that engineers attack the largest loss measured on the job timeline rather than whichever subsystem they inspected most recently, while the mechanisms that shorten each loss remain with their owning chapters.

9.5.1 What runtime goodput counts

Holding a complete gang for an hour does not guarantee an hour of retained training steps: the job can spend much of that hour launching processes, waiting on a starved data pipeline, pausing for synchronous checkpoint commits, or re-executing steps erased by a node failure. Runtime goodput measures the fraction of all-allocated time that survives those interruptions as a ratio of two durations over the same calendar window. The denominator is the job's held time from the gang-complete timestamp to release, taken directly from scheduler logs (section 9.4.3). The numerator is productive time: the count of retained steps multiplied by the job's fault-free step duration. A step counts as retained once a committed checkpoint records its state update (or when the job finishes and commits its final weights). The fault-free step duration is the job's steady-state step time with input batches ready and no recovery in progress on that same allocation. It is the mean input-ready step time over every retained step outside post-restart warmup, so intermittent slowdown episodes stay in the baseline and are charged to program goodput.

Normalizing the numerator to the job's own fault-free step duration keeps the runtime boundary cleanly separated from the program boundary. Because the fault-free step duration already includes exposed collectives, pipeline bubbles, memory bandwidth stalls, and persistent stragglers, those per-step inefficiencies pass through the runtime boundary unchanged and are evaluated in program goodput. Steps immediately following a restart are handled differently. While instruction caches warm up, just-in-time kernels compile, and data-loader queues refill, any excess step duration above the steady-state fault-free baseline belongs to the recovery event that triggered the restart and is charged to runtime goodput. The fault-free baseline duration of each completed step, by contrast, counts as retained progress. Consequently, a degraded node that slows every step across a multi-day run registers as a program goodput loss, whereas a temporary slowdown confined to the warmup minutes after a restart registers as a runtime goodput loss.

Every other interval of held time that fails to produce retained steps is a runtime loss, ordered chronologically across a job's lifecycle:

- *Startup* spans the interval from gang completion to the first training step while processes launch, libraries load, collective communicators initialize, checkpoint weights load, and compilers warm up. Because scheduler logs timestamp gang completion and the job timeline timestamps the first step, startup is measured directly between those two records—which is why section 9.4.3 requires the scheduler, rather than the first step, to mark allocation completion.
- *Input waits* capture the time a retained step spends blocked on the data pipeline for its next batch, carved out of step duration per section 9.2.3.
- *Synchronous checkpoint pauses* capture the intervals during which training halts while accelerators or hosts flush state to persistent storage.
- *Failure recovery and rework* cover both the downtime from a fault through detection, relaunch, state reload, and warmup, and the held time spent earlier on post-checkpoint steps that the failure erased and the job had to repeat.
- *Forced preemption and planned reconfiguration* charge the post-checkpoint work discarded by an eviction, the subsequent reload after readmission, and the warmup before steady-state step times return, along with planned maintenance drains or elastic resizes. All of these belong to runtime goodput because the job held a complete allocation while the discarded or non-step hours were spent.

Three adjacent delays fall outside runtime goodput. Time a preempted job spends waiting in the scheduler queue before readmission is not held time and affects only time to result (section 9.4.1). Background bandwidth contention from asynchronous checkpoint writes and the steady-state execution speed of compiled kernels both lengthen retained steps and belong to program goodput, whereas ahead-of-time compilation before the first step belongs to runtime startup. Keeping cold-start compilation and warmup in runtime goodput matches the measurement contract of table 9.2, which excludes warmup trials from the steady-state step rate used for program goodput.

Finally, the measurement window for runtime goodput must span many interruption cycles (section 9.3.1) and requires a clean closing rule: steps executed after the last checkpoint committed inside the window are neither durably retained nor lost yet. Closing the measurement window at a checkpoint commit avoids that ambiguity; if a window closes between commits, those trailing steps count as retained only once the next checkpoint succeeds.

9.5.2 Composing runtime loss from its owners

A team whose job loses a tenth of its held time cannot tell from the aggregate runtime goodput alone whether to add storage bandwidth for input stalls, shorten checkpoint writes, or negotiate a tighter preemption budget. Because every non-productive interval of held time falls into exactly one of the five runtime-loss categories, those categories partition the gap between held time and productive time without overlap. Dividing each category's duration by total held time expresses it as a fractional loss share, and the five shares add directly. Earlier chapters already derived each individual share:

1. The *startup share* is the launch and compilation duration T_{startup} from gang completion to first step, divided by the held duration T_{wall} of that allocation.
2. The *input-stall share* is the data stall ratio f_{stall} of equation 4.4, which section 4.4 derives from the bottleneck stage of the storage pipeline.
3. The *checkpoint share* is the Young-Daly checkpoint waste $W(T_{\text{ckpt}})$ of equation 7.8 evaluated at the job's checkpoint interval T_{ckpt}, combining synchronous write pauses with expected post-checkpoint rework after failures.
4. The *failure-recovery share* is the per-fault recovery duration T_{recovery} from section 7.9.2 (detection, relaunch, checkpoint load, and warmup) divided by the cluster's mean time between failures $\text{MTBF}_{\text{system}}$.
5. The *preemption share* is the forced-preemption rate λ_{preempt} multiplied by the per-eviction overhead T_{preempt} from section 8.8.1.

Equation 9.4 sums the five shares to give total runtime goodput loss:

$$1-\eta_{\text{runtime}} \approx \frac{T_{\text{startup}}}{T_{\text{wall}}} + f_{\text{stall}} + W(T_{\text{ckpt}}) + \frac{T_{\text{recovery}}}{\text{MTBF}_{\text{system}}} + \lambda_{\text{preempt}}\, T_{\text{preempt}} \tag{9.4}$$

Writing runtime loss as an explicit sum turns optimization into a quantitative triage problem for a specific workload: which of the five terms dominates *this* job at *this* scale under *this* cluster policy. Each term carries the modeling assumptions of its owning chapter, so whenever a workload violates one of those assumptions, engineers refine that specific term using the owning chapter's full model while keeping the additive boundary accounting intact.

↩ The checkpoint waste equation, its optimal interval, and the four-term recovery budget are established in chapter 7.

Equation 9.4 is a first-order sum because the individual models normalize their losses against slightly different baselines—step time for data stalls, fault-free compute time for Young-Daly waste, and arrival rates times event durations for recovery and preemption. When every individual loss share is small, those baseline differences alter the sum only to second order (fractions of a percent). Whenever a single term grows large enough to consume a major fraction of held time, the first-order approximation has already served its diagnostic purpose, and engineers measure the exact shares directly from the job timeline.

Summing the five terms is valid only when no lost minute is counted in two terms at once. Four boundary junctions require care so that the owning chapters' formulas do not overlap:

- *Rework is counted once per interruption cause.* Because the checkpoint waste term $W(T_{\text{ckpt}})$ already includes the expected steps lost after a hardware failure, the failure-recovery term

$T_{\text{recovery}}/\text{MTBF}_{\text{system}}$ includes only detection, relaunch, reload, and warmup. Conversely, steps erased by a scheduler eviction belong to T_{preempt} and must not be folded into $\text{MTBF}_{\text{system}}$ unless the two terms are solved jointly.

- *Warmup is charged to the event that triggered it.* Both the failure-recovery budget T_{recovery} and the preemption tax T_{preempt} already include post-restart warmup, so the initial startup term $T_{\text{startup}}/T_{\text{wall}}$ counts only cold launches that were not triggered by an in-run failure or preemption.
- *The data stall ratio f_{stall} is measured in steady state.* Pipeline refill immediately after a restart is part of post-restart warmup; averaging f_{stall} across restart warmup windows would charge that refill delay twice.
- *Only synchronous checkpoint pauses enter $W(T_{ckpt})$.* When checkpoints are written asynchronously in the background, any memory or fabric contention they inflict on running steps lengthens retained step time and is charged to program goodput instead.

At the 10,000-GPU scale analyzed in section 7.8.1, checkpoint waste (about 5.9 percent) and failure recovery (about 2.8 percent) together consume nearly a tenth of held time. Both shares are evaluated here at the full component failure rate that governs the Young-Daly interval, so this recovery share exceeds the GPU-only recovery figure of section 7.9.2. Both failure-driven terms grow with device count, whereas startup, input stalls, and preemption do not. Consequently, jobs of different sizes lose runtime goodput to entirely different mechanisms. Because our running example is much smaller than that frontier reference, its loss terms rank in a different order.

Napkin Math 9.3: Runtime loss of the Llama 2 70B job

Problem: The Llama 2 70B training job holds 128 GPUs on a shared fleet for a week, under a per-job budget of 1 forced preemption a day. With each owner's model applied to the job's own configuration, what runtime goodput does it keep, and which term costs it the most?

Variables:

- Mean time between failures $\text{MTBF}_{\text{system}}$ is about 302 hours, from the component failure rates of section 7.8.1.
- The resumable checkpoint takes about 9.8 seconds to write at the 100 GB/s of the recovery budget in section 7.9.2, the job's tenth of a parallel file system shared by ten jobs (section 4.6), so the Young-Daly interval T_{ckpt} is about 77 minutes.
- Recovery T_{recovery} is about 6.2 minutes per failure, and startup T_{startup} is about 5.2 minutes per launch, the same steps without detection.
- The data stall ratio f_{stall} is 7.7 percent of held time, the snapshot's 5 percent of in-service capacity divided by its scheduling goodput of 0.65 (section 9.4.3).
- The preemption overhead T_{preempt} is about 20.2 minutes per forced event, the 15 minutes of lost work that section 8.8.1 assumes plus one startup.

Math (one line per term of equation 9.4):

- Startup: 5.2 minutes $\div$ 7 days $\approx$ 0.05 percent.
- Input waits: 7.7 percent.
- Checkpoint: 9.8 seconds $\div$ 77 minutes for writes, plus 77 minutes $\div$ (2 $\times$ 302 hours) for rework, $\approx$ 0.2 percent + 0.2 percent $\approx$ 0.4 percent.
- Recovery: 6.2 minutes $\div$ 302 hours $\approx$ 0.03 percent.
- Preemption: 1 per day $\times$ 20.2 minutes $\approx$ 1.4 percent.

Result: The five terms add to a runtime loss of about 10 percent, a runtime goodput of about 0.90. Input waits and preemption account for 9.1 percent of it, and checkpoint waste and recovery together for under half a percent.

Systems insight: On a mid-sized job that shares its fleet, the terms that grow with device count are the smallest, so the lever that reaches the most held time belongs to the input pipeline and the preemption budget rather than to the reliability layer.

Because the overall runtime goodput is already close to what a well-run job achieves, the engineering lesson lies in how the loss is divided. A team consulting only chapter 7 might invest in faster checkpoint writes and warm restarts, yet on this 128-GPU job those levers can recover at most the 0.4 percent and 0.03 percent shares assigned to them by the composition. Meanwhile, the input pipeline sized in section 4.4 governs roughly 17× as much held time. Without a preemption cap, the 12 forced preemptions per day that section 8.8.1 models for a low-priority run would consume 16.8 percent of held time—exceeding every other term and breaking the small-loss regime of the first-order sum.

Both job size and job duration reorder these terms. As device count grows, checkpoint waste and failure recovery expand until their combined loss overtakes input waits near 13,000 GPUs (figure 9.5). Similarly, the same 5.2 minutes startup cost that is negligible over a week consumes about 9 percent of an evaluation job that holds its allocation for only an hour. The composition does not pick a universal winner among these levers; rather, it shows each subsystem owner how much held time their lever can actually reach—a comparison no single subsystem can make in isolation.

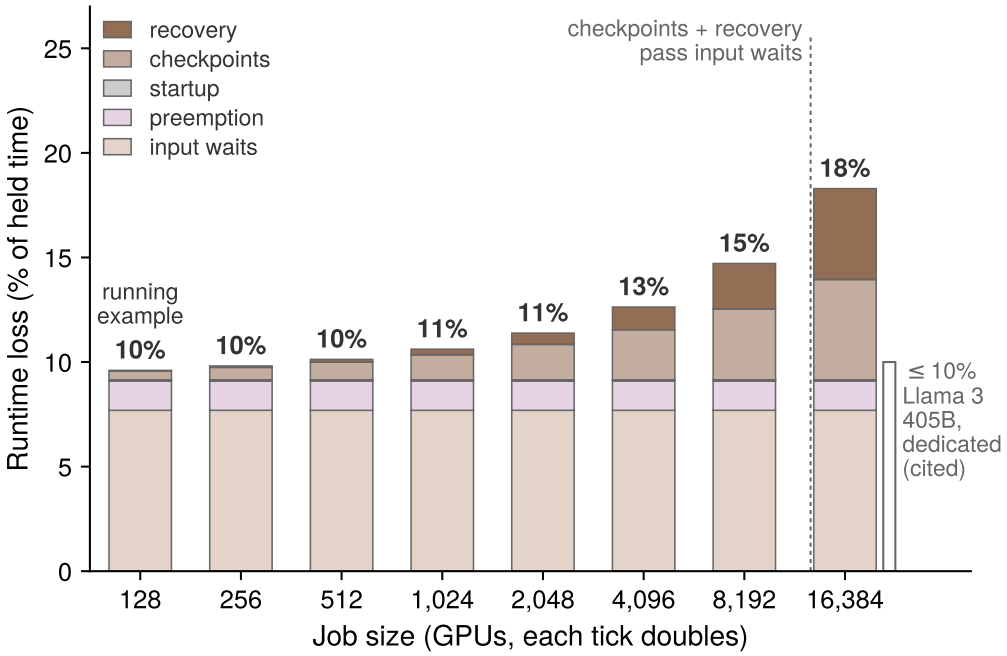

Figure 9.5: Runtime Loss versus Job Size: Each bar recomputes the five terms of equation 9.4 for the running example at one job size, from 128 to 16,384 GPUs, holding the 70B checkpoint, the recovery budget, the data stall ratio, the preemption budget, and the week of held time at the running example's values and recomputing the Young-Daly interval at each size. Input waits and preemption stay flat and dominate small jobs, while checkpoint waste and recovery grow with device count, pass input waits near 13,000 GPUs (dotted line), and lift the loss from 10 percent to 18 percent, above the loss of at most 10 percent implied by the effective training time Llama 3 405B reported on its dedicated cluster (hollow bar) (Dubey et al. 2024).

Composing the terms also exposes a cross-layer coupling that no single owner sees alone. Because a forced preemption discards all work since the last committed checkpoint just as a hardware failure does, it acts as an additional interruption when choosing the checkpoint cadence. Our running example's 77 minutes interval was derived from hardware failures alone; counting its budgeted preemptions as interruptions in equation 7.7 shortens the optimal interval to about 21 minutes. Although the checkpoint cadence belongs to section 7.8.1 and the preemption budget belongs to

section 8.8.1, composing runtime goodput reveals that the two parameters are coupled parts of a single decision.[5]

While equation 9.4 describes an individual job, a pool's aggregate runtime goodput is also lowered by idle reserved blocks and by zombie allocations that continue to hold devices after execution has stalled. Because section 9.4.2 charges both intervals to runtime goodput, they appear only when the pool's total held time is compared against the sum of its active jobs' timelines.

9.5.3 Evidence, size, and the serving replica

A job that hangs inside a collective past its detection timeout, warms up much more slowly than its recovery budget assumed, or compiles longer after each software release loses held time that no analytical model predicts. Measuring runtime goodput directly from the job timeline—the second rung of section 9.3.2—captures both the modeled terms and any unmodeled residual because the timeline's events map directly onto the loss categories. Startup runs from the scheduler's gang-complete event to the first step; input stalls come from the per-step input-wait field of each retained step; and checkpoint pauses appear as the intervals at each commit. Every failure or preemption contributes its detection, relaunch, and warmup intervals together with the steps it forced the job to repeat, while excluding any queue wait after an eviction. Finally, the retained-checkpoint markers establish which steps count toward progress, and the steady-state step times fix the fault-free rate.

Under this workflow, equation 9.4 serves as an analytical check on the job timeline rather than a substitute for it. When prediction and measurement agree, the largest modeled term identifies the subsystem owner to call. When the measured loss exceeds the prediction, the unexplained residual represents held time lost to unmodeled hangs, slow warmups, or creeping initialization overheads. A residual that widens across many jobs after a deployment is precisely the regression caught by the runtime row of table 9.4.

The Llama 3 405B training run provides the benchmark for a well-run frontier job. Meta reported more than 90 percent effective training time, corresponding to a runtime goodput of at least 0.90 (Dubey et al. 2024)—a demanding standard given the hardware failure cadence analyzed in section 7.1.

Although our running example's runtime goodput of 0.90 matches that frontier figure to two decimal places, the agreement masks a difference in mechanism. Running on a dedicated cluster across 128× as many devices, the Llama 3 job had to fight for its runtime goodput against failure-driven checkpoint and recovery overheads. Our 128-GPU job loses almost nothing to hardware failures and nearly all of its lost tenth to input stalls and preemption. Because the same aggregate factor can point to entirely different owners, a fleet report must record the constituent loss terms alongside their sum.

By contrast, a standard benchmark run is short, dedicated, and timed only after initialization completes,[6] so of the five runtime loss terms it experiences at most the steady-state input stall (perspective 9.1). Any capacity plan extrapolated from a benchmark rate must therefore restore the remaining terms of this composition.

For a serving replica, the same boundary separates held time from ready time, so its runtime goodput is the ratio of ready time to held time in the replica's lifecycle log. Provisioning precedes the complete allocation and belongs to scheduling goodput, whereas the weight fetch, model load, and warmup that follow—requiring about 5.7 minutes in the cold-start budget of section 10.13—are runtime losses. How heavily those minutes weigh on runtime goodput depends on the replica's lifetime: a replica recycled once a day spends about 0.4 percent of its held time becoming ready, whereas a replica spun up for a single burst hour loses about 9 percent. Maintaining a warm pool of loaded, idle replicas does not eliminate this cost; it shifts the idle capacity from runtime goodput into program goodput as serving headroom.

Whether measured as a ready replica or a retained training step, crossing this boundary means that held time has become retained work. How fast that retained work executes relative to its hardware ideal is the domain of the third factor: program goodput.

5 **Effective training time ratio (ETTR)**: Kokolis et al. (2025) define a job's ETTR as productive runtime over wallclock time, including time spent queued, and derive its expected value from the failure rate, checkpoint interval, checkpoint write cost, restart overhead, and queue time; they use it to estimate how much large jobs on Meta's research clusters lose to failures. It is the closest published counterpart to this composition for the failure-driven terms. Because its denominator includes queueing, ETTR also spans part of scheduling goodput, so an ETTR and a runtime goodput for the same job are not directly comparable.

6 **MLPerf Training timing**: The rules as first published start the clock when the system first touches training or validation data, leave system initialization untimed, and exclude up to 20 minutes of model creation as well as one-time data reformatting (Mattson et al. 2020). Startup, compilation, and data preparation therefore never enter a benchmark's time to train, although every production launch pays them in the startup term.

↪ Cold-start latency, its decomposition into provisioning, loading, and warmup, and the pre-warming that hides it are developed in section 10.13.

Checkpoint 9.1: Charging losses to factors

Every lost interval is charged once, at the first boundary it fails to cross, and each factor is read from the record that can see its boundary.

- ☐ **Which factor**: For a gang that holds half its devices while it waits, a zombie job, an input wait inside a retained step, and a straggling rank, can you name the factor each lowers and the layer whose lever removes it?
- ☐ **Which record**: A job retains fewer steps than planned. Can you explain why the scheduler records and the job timeline are read before any step trace, and what each can and cannot establish?
- ☐ **Which window**: Why can a one-day measurement report a runtime goodput near one for a job that loses a tenth of its held time over a month?
- ☐ **Composing runtime loss**: Given a data stall ratio, a checkpoint interval, a failure rate with its recovery time, and a preemption budget, can you add the runtime loss terms and name the interval that must not be counted twice?

9.6 Program Goodput

Even when a job holds a complete allocation and retains every step it executes, it can still waste most of its capacity if each retained step runs far below the speed its hardware allows. Within the held time of our Llama 2 70B training job, program goodput is the single largest source of loss—even though scheduling losses consume a larger share of the shared pool as a whole (section 9.8.1). Because the job's retained steps run at a fleet MFU of 48 percent, its program goodput is about 0.48, meaning that 52 percent of the time spent inside retained steps produces no progress at the ideal rate.

Program goodput is also the only factor whose losses this chapter both diagnoses and recovers directly by changing how the step executes. Diagnosing a program goodput loss requires two complementary instruments: the iron law and the Fleet Law identify which analytical term of a step a loss lengthens, while a hierarchical step trace—read from the fleet metric down to individual hardware counters—shows which term is actually exposed on the critical path.

↩ Partitioning layouts and their communication patterns, which set most of a large job's step structure, are established in chapter 6.

9.6.1 The iron law and the Fleet Law terms

Speeding up a kernel on a single accelerator often yields much less benefit—or none at all—once a job scales across a fleet, because the step also spends time moving data across memory hierarchies, waiting at synchronization barriers, and exchanging tensors over the network fabric. Separating what a local kernel change can reach from what requires a fleet-level change starts with the single-machine iron law (equation 9.5). The iron law decomposes the local execution of each device into the data volume crossing a chosen memory boundary, the arithmetic operations performed, and the latency overhead that neither term explains:

$$T = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}} \tag{9.5}$$

Here D_{vol} is the data volume crossing the chosen memory boundary, BW is that boundary's bandwidth, O is the operation count, and $R_{\text{peak}} \cdot \eta_{\text{hw}}$ is the realized compute rate. At fleet scale, the overhead term L_{lat} splits into two parts with distinct owners: host dispatch (the kernel launches that keep each accelerator fed) belongs to single-device execution, whereas exposed communication and synchronization exist only because the step spans multiple devices. Section A.2 reviews these terms and their overlapped form.

The Fleet Law (equation 1.3) models a distributed step as each device's share of local work plus communication time plus synchronization time, minus the portion of communication and synchronization hidden behind computation. The single-machine iron law connects to the fleet through these terms. The data-movement and compute terms, together with the host-dispatch portion of L_{lat}, constitute local execution, whose unsharded reference time is T_{compute} and whose per-device share across N devices is T_{compute}/N. Meanwhile, the communication component of

L_{lat} becomes the Fleet Law's $T_{\text{comm}}(N)$, the synchronization component becomes $T_{\text{sync}}(N)$, and any communication or synchronization hidden behind useful computation is credited back through T_{overlap}. Profiling one device of an already-partitioned job measures the per-device share of local execution directly, because partitioning has already divided the local workload, so the measured time must not be divided by N a second time. Throughout the remainder of this chapter, "overhead" refers to host dispatch unless communication is named explicitly. A local kernel optimization therefore improves fleet performance only through the specific term it shortens, and it shortens the overall step only by the portion of that term that was exposed on the critical path.

Every program goodput loss cataloged in table 9.1 lengthens one of these Fleet Law terms. Memory stalls, sub-peak arithmetic, and slow host dispatch all lengthen local execution, whereas exposed collectives lengthen $T_{\text{comm}}(N) - T_{\text{overlap}}$. Stragglers and pipeline bubbles both manifest in $T_{\text{sync}}(N)$ as healthy devices wait at a barrier or stage boundary for upstream work to finish, whereas activation recomputation inflates O beyond the work required by the model itself. Each mechanism leaves a distinct signature in a step trace and answers to its own engineering lever (table 9.6).

Table 9.6: Program Goodput Losses by Fleet Law Term: Each loss inside a retained step lengthens one term of the iron law or the Fleet Law, and a step trace shows it as a characteristic interval. Only barrier and stage waits of the Fleet Law's coordination term are charged here.

Program goodput loss	Term it lengthens	Evidence in a step trace	Lever owner
Memory stalls	D_{vol}/BW in local execution	Kernels near the bandwidth ceiling; achieved HBM bandwidth close to peak while arithmetic is low	This chapter (fusion, tiling, compilation)
Arithmetic below peak	$O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$ in local execution	Kernels far below both roofline ceilings; low occupancy or poor instruction mix	This chapter (kernel selection, compilation)
Recomputation	Added O beyond the model's required work	Hardware FLOPs utilization above MFU	chapter 6 (memory feasibility)
Host dispatch	Dispatch part of L_{lat} in local execution	Idle gaps between short kernels; long host-side launch intervals	This chapter (graph replay, compilation)
Exposed communication	$T_{\text{comm}}(N) - T_{\text{overlap}}$	Collective intervals with no concurrent computation on any rank	This chapter and chapter 5
Stragglers	$T_{\text{sync}}(N)$	All ranks but one waiting at the same barrier	chapter 7 (replacement)
Pipeline bubbles	$T_{\text{sync}}(N)$	Stages idle at the start and end of each step	chapter 6 (schedule)

Although the coordination term of the Fleet Law also encompasses checkpoint writes, failure detection, and failure recovery, those intervals occur *between* retained steps and are charged to runtime goodput (section 9.5). The Fleet Law's coordination term therefore straddles two goodput factors: only the barrier and pipeline-stage waits *inside* a retained step lower program goodput. Put simply, the Fleet Law term identifies the physical mechanism and its engineering lever, while the goodput boundary identifies which factor pays the cost.

This mapping also bounds how much a single-device speedup can help at scale. Because a kernel optimization shortens only T_{compute}/N, at a fixed global batch its share of the total step shrinks as N grows and T_{compute} is divided across more devices, while communication and synchronization do not shrink in proportion. Under weak scaling, as in the running example, each device's share of local work stays constant while communication and synchronization grow, so the kernel's share shrinks in that regime too. Consequently, the very same kernel speedup delivers a smaller end-to-end gain on a larger job. Optimizations can also shift cost from one term to another rather than eliminating it: recomputing activations to free HBM capacity adds floating-point operations to O, illustrating that the displacement of overhead (principle 14) operates within a single step just as it does across the C^3 dimensions.

The serving replica applies the exact same decomposition to a single decode step. Because each of its eight devices holds one shard of the 70B model's weights, a small-batch decode step must stream its entire weight shard from HBM to perform only a handful of operations per weight, making the data-movement term D_{vol}/BW dominant. Meanwhile, the tensor-parallel activation exchange in

every transformer layer forms its communication term. The replica's program goodput is therefore governed by how closely its decode step approaches the minimum time required to stream its weights—a relationship that the roofline model makes concrete in section 9.6.3.

In the replica's small-batch decode step, the data-movement term dominates the iron law.

> **Checkpoint 9.2: The iron law and the Fleet Law terms**
>
> Program goodput falls only through the terms of a step, and a change helps only through the term it shortens.
>
> - ☐ **Iron-law estimate**: For a step that is 60 percent data movement, 30 percent compute, and 10 percent overhead, can you estimate which single change buys the most wall-clock time, and what the iron law predicts once that dominant term is halved?
> - ☐ **Fleet Law term**: Can you name the Fleet Law term a fused kernel shortens, and explain why its fleet gain shrinks as N grows?
> - ☐ **Barrier waits**: Can you explain why a rank that waits at a barrier for a slow peer is charged to program goodput, while a job that waits for a failed node to be replaced is charged to runtime goodput?

9.6.2 Exposed communication at fleet scale

A job distributed across many nodes must exchange data over the inter-node network fabric, which is slower than intra-node links, shared with competing jobs, and traversed by collectives whose volume and latency grow with participant count. Whatever portion of that exchange the execution schedule cannot hide behind useful computation lengthens every retained step. That exposed communication interval is what separates a multi-node fleet job from the same workload on a single node, and because the job holds a complete allocation and retains every slowed step, the loss belongs entirely to program goodput.

↩ How much communication a schedule can hide, and why the hidden fraction falls as collectives grow relative to the compute window, are established in section 5.10.4.

Our running example isolates this multi-node penalty in a single comparison, with constructed step times standing in for measured ones. Because its single-node baseline and its 128-GPU fleet configuration train the same model with the same per-device workload, the ratio of their MFUs estimates what scaling across the fabric costs the program. It is an estimate rather than an exact measure because the two configurations can also differ in other ways, such as the node baseline's offload of optimizer state to host memory.

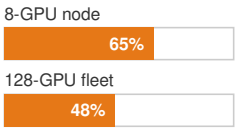

Fleet MFU falls below the single-node baseline.

> **Napkin Math 9.4: From one node to 128 GPUs**
>
> **Problem**: The 70B parameter training job of the running example achieves a node baseline of 65 percent MFU on one 8 GPUs node, with optimizer state offloaded to host memory because the full training state of a 70B model exceeds one node's HBM. Scaled to 128 GPUs with the same work on each device, its constructed microbatch-scale step time rises from 180 ms to 245 ms. What program goodput does the job keep at fleet scale, and how much of the node's efficiency does scale cost?
>
> **Math**: Both MFU figures apply equation 9.3, the useful FLOPs of a step divided by the FLOPs the devices could execute at peak in that step's wall time.
>
> Step 1: Fleet MFU. A step performs 14.9 PFLOP of useful work, while 128 GPUs $\times$ 989 TFLOP/s per GPU $\times$ 245 ms could execute 31 PFLOP at peak, so fleet MFU is 14.9 PFLOP $\div$ 31 PFLOP $\approx$ 48 percent. The same substitution on the node gives 65 percent.
>
> Step 2: Efficiency relative to the node. Because each device does the same work in both configurations, the ratio of the two MFUs is the ratio of the two step times, 180 ms $\div$ 245 ms $\approx$ 74 percent.
>
> Step 3: Time beyond the node. Each fleet step spends 180 ms on the node's work plus 65 ms that the node step did not, about 27 percent of the fleet step.
>
> **Result**: The job keeps program goodput of about 0.48 and 74 percent of the node's efficiency, losing 26 percent of it to scale.

Systems insight: The 65 ms per step beyond the node baseline bounds what exposed communication and synchronization can claim, but it does not by itself isolate the fabric, since offload, kernel shapes, and barrier waits can also change between the two configurations. The bound becomes an attribution only when a step trace shows which intervals fill the added time.

The ratio computed in the notebook is the scaling efficiency of equation 6.1—referenced here to a single node rather than to a single device as in section 6.4.2—and it completes the metric reconciliation begun in section 9.2.2. A job's fleet MFU equals its single-node baseline MFU multiplied by its scaling efficiency from that node (65 percent $\times$ 0.74 $\approx$ 48 percent), and that product, evaluated over retained steps, is its program goodput. An engineering report may quote the fleet MFU directly, or it may quote the single-node MFU alongside the scaling efficiency, but it must never multiply the fleet MFU by the scaling efficiency a second time, which would double-count the multi-node penalty. Because the single-node baseline already includes intra-node NVLink collectives and host-offload traffic, factoring fleet MFU at the node boundary bounds the cross-node scaling cost against a baseline that engineers can easily reproduce.

In the language of the Fleet Law, the additional 65 ms per step reflects growth in $T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}}$ alongside any configuration difference in local execution, and only a step trace can separate those contributions. In the constructed diagnosis that section 9.3.2 walked through, the trace showed no rank consistently lagging its peers but did show exposed cross-node synchronization and collective waits on every microbatch step. Had the trace instead shown fifteen nodes waiting at a barrier for one slow node, the very same step-time increase would have belonged to $T_{\text{sync}}(N)$—yielding the same program goodput, but pointing to node replacement rather than to the communication terms as the place to act.

The Llama 3 405B training run provides a production reference point: it sustained an MFU of about 41 percent across 16,384 GPUs[7] under a parallelism layout tuned by the team that built the cluster (Dubey et al. 2024). The 128-GPU running example's slightly higher program goodput is a constructed value, not a measured comparison. It was set above Llama 3's figure to reflect the expectation that a job spanning far fewer devices spends a smaller fraction of each step in cross-node collectives, and it says nothing about the quality of either run's kernels. Both runs surrender roughly half of their retained step time relative to the peak arithmetic ceiling; the task of performance diagnosis is to determine which specific Fleet Law terms hold that lost half.

[7] **MFU across Llama 3 configurations**: The report's scaling table gives 43 percent BF16 MFU on 8,192 GPUs and 41 percent on 16,384 GPUs, with the same tensor and pipeline layout and twice the data parallelism (Dubey et al. 2024). The report attributes the drop to the smaller batch per data-parallel group needed to hold the global batch constant, so doubling the devices cost program goodput with no change to the model or its kernels.

9.6.3 Drilling into the step

Inspecting hardware counters on an arbitrary kernel before knowing which rank or stream delayed the step wastes engineering effort on operations that may not even lie on the critical path. Once a loss has been localized to program goodput, the step trace becomes the primary evidence, and it must be read from the top down. Diagnosis descends through four hierarchical levels, and the lower three expand the step-trace rung of section 9.3.2 (figure 9.6):

1. **Job and fleet metrics** (step time against plan, MFU of retained steps, and per-rank step timers) detect that a program goodput loss exists.
2. **Framework and distributed traces** correlate step time with model operations across ranks, revealing whether a lagging rank or an exposed collective is lengthening the step.
3. **Kernel timelines** expose host launch gaps and the stream-level overlap between computation and communication on a single rank.
4. **Hardware performance counters** measure achieved memory bandwidth, SM occupancy, and instruction mix for an individual kernel.

Each lower level is collected only for the specific ranks and time windows flagged by the level above it. Section A.8 maps the single-device profiling tools used at each rung.

The constructed diagnosis of our running example descends three of the four levels. First, distributed traces across a data-parallel group show that no rank lags consistently behind its peers, which rules out a hardware straggler, and they reveal collective intervals that widen with device count. Next, a kernel timeline on a representative rank shows those collectives executing on the communication stream with no concurrent kernels on the compute stream during the latter part

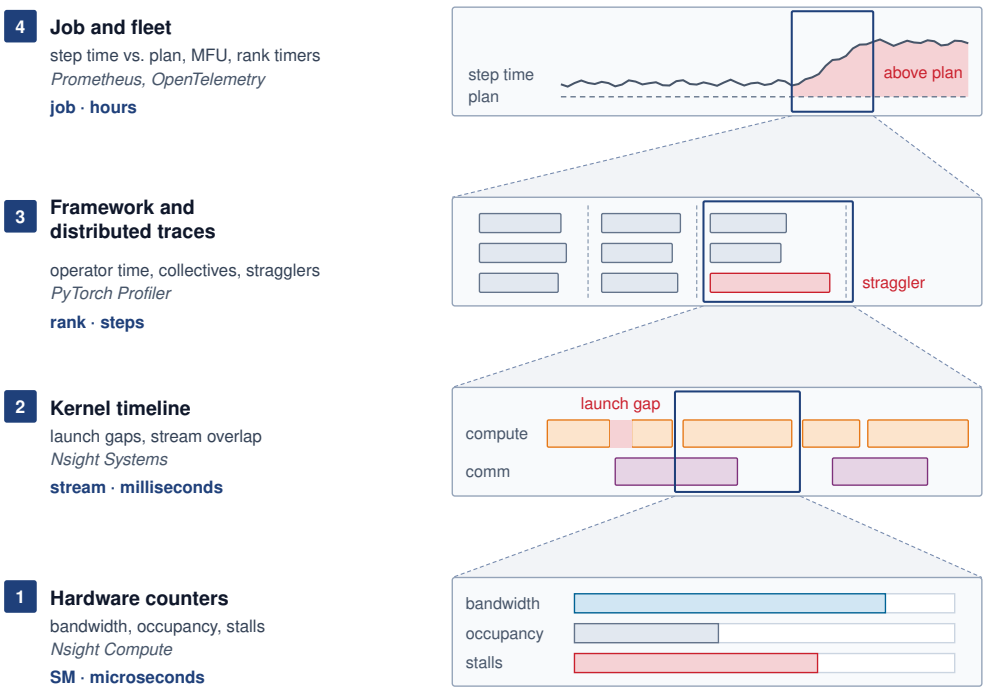

Figure 9.6: The Profiling Hierarchy: Diagnosis of a program goodput loss descends from the fleet measure that detected it to the hardware root cause. Each rung names the evidence that level exposes and a representative tool; the navy label under each tool gives the scope and time scale one sample covers, and each highlighted window expands into the trace below it. Each lower level trades breadth for the detail needed to isolate a root cause, and the three lower levels together are the step-trace rung of factor localization.

of the backward pass. That timeline establishes that those collectives are exposed, but not why. Separating a poor collective schedule from a transfer too long for the backward compute window requires comparing the collective's transfer time with that window (section 5.10.4). For the job's data-axis gradient reduction the comparison is decisive, because that transfer takes longer than the entire microbatch step. No schedule can hide it within one step, so the job runs it once per optimizer step after accumulating gradients over many microbatches (section 9.7.2). The exposed intervals on every microbatch step are, by assumption of the constructed example, per-microbatch synchronization across the nodes and the collective waits it causes, not a transfer the example sizes. Hardware counters, the lowest level, are unnecessary because the compute kernels flanking the exposed collectives already run at their expected rates, and the descent stops as soon as the responsible Fleet Law term is identified.

Systems Perspective 9.2: The cost of tracing a step

Tracing adds work to the step it measures. Recording every kernel launch, memory allocation, and collective consumes host cycles and memory and can add synchronization, so a traced step can run slower, or with a different overlap between communication and computation, than the untraced steps it is meant to explain. The perturbation grows with the events and options collected. A trace is therefore evidence about the structure of a step, while the step's duration is measured from the untraced job timeline, and comparing an instrumented with an uninstrumented window of the same job shows how much the trace changed.

The window must also describe the job's steady state. A profiler's own warmup steps let the profiler start up before it records, but they do not ensure that the application has finished

compiling, allocating memory, and initializing its devices, so the window is chosen from step times that have stopped changing. Startup that falls inside the trace belongs to runtime goodput and is measured from the timeline, not from the trace.

A first-order bound on program goodput needs no trace at all. The diagnostic from section 8.12.3 compares total step duration against compute-only step time (the forward and backward passes alone). Because ideal hardware step time cannot exceed compute-only time, the ratio of compute-only time to actual step time is an immediate upper bound on program goodput.

Systematic triage of a slow step begins outside the accelerator kernels. If a step misses its duration target while accelerator utilization is low, the bottleneck lies in input stalls, host preprocessing, or exposed communication—because optimizing a kernel cannot speed up a device that is waiting for work to arrive. Only when the accelerator is actively executing kernels does diagnosis turn to whether those kernels are bound by arithmetic throughput or memory bandwidth, evaluated against the roofline's ridge point (the operational intensity where the memory-bandwidth slope meets the compute ceiling). Organized in order of application, these checks form a diagnostic decision tree (figure 9.7) whose initial branches exit program goodput altogether. A device waiting for its next batch or for host-side preprocessing is losing runtime goodput at the input boundary, which returns the investigation to section 9.5.

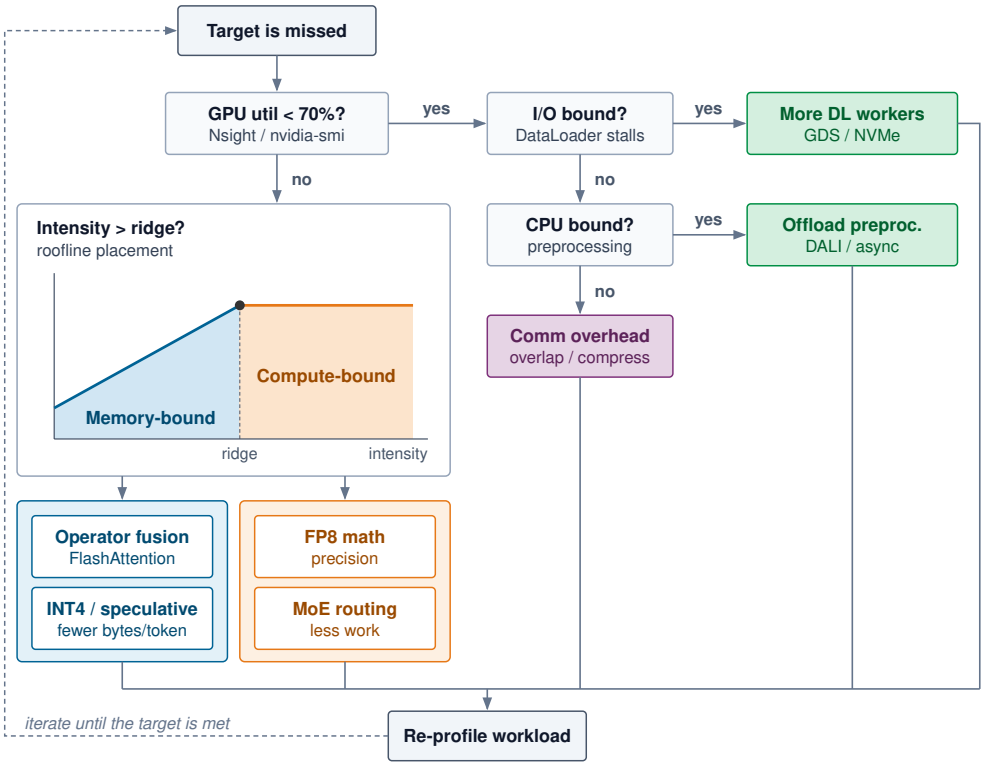

Figure 9.7: Iron Law Diagnostic Flowchart: A triage flow for a job or replica whose retained steps miss their step-time target. Low accelerator activity sends the diagnosis right, toward waits, and high activity sends it down to the roofline. The I/O and CPU branches identify input waits and host preprocessing, which are runtime goodput losses at the input boundary, so a yes on either sends the reader back to section 9.5 rather than to the program. The communication-overhead leaf is exposed communication, a program goodput loss in the Fleet Law's communication term. Only a busy device reaches the roofline panel, whose memory-bound and compute-bound leaves are candidate levers to test, not prescriptions. Fusion and tiling are developed in section 9.7, serving precision and speculative decoding in chapter 10, and expert routing in section 6.5.5. Every path returns to re-profiling.

Every path in the flowchart concludes by re-profiling the workload, because an optimization counts only if the measured step becomes shorter under the contract of section 9.3.1. Removing one exposed interval frequently uncovers the next, shifting the critical path and restarting the diagnostic cycle.

Once the flowchart confirms that an accelerator is busy, its kernels are evaluated against the roofline model. A kernel with arithmetic intensity I can achieve at most $\min(R_{\text{peak}}, I \times \text{BW})$, and the ridge point $I_{\text{ridge}} = R_{\text{peak}}/\text{BW}$ divides the memory-bound slope of that envelope from its compute-bound plateau (Williams et al. 2009). Whereas section 2.3.2 introduced the Roofline Model from hardware specifications and algorithmic bounds, performance engineering places each kernel on the roofline using the operations and memory traffic measured during the actual step that lost program goodput.

Definition 9.4: Measured arithmetic intensity

Measured arithmetic intensity is the ratio of the floating-point operations a kernel actually executes to the bytes observed crossing a named memory boundary, such as HBM, during that execution (FLOP per byte).

1. **Significance**: It places a profiled kernel on the same roofline as the bandwidth (BW) and compute (R_{peak}) ceilings, so its distance below the algorithmic estimate exposes traffic the algorithm did not require.
2. **Distinction**: Unlike *algorithmic intensity*, which counts the data movement an idealized algorithm requires, measured intensity includes tiling, cache behavior, fusion, padding, and redundant transfers.
3. **Common pitfall**: A roofline point is interpretable only when its FLOP convention and memory boundary match the ceilings on the same plot.

Placing a profiled kernel on the roofline proceeds in three steps. First, hardware counters for the kernel report its achieved FLOP rate and its achieved bandwidth across a named memory boundary. Second, dividing the achieved FLOP rate by the achieved bandwidth yields the kernel's measured arithmetic intensity, which reveals excess memory traffic whenever it falls below the algorithmic estimate. Third, comparing the achieved FLOP rate against the roofline ceiling at that measured intensity classifies the kernel into one of three regimes:

1. **Far below both ceilings.** The kernel suffers from an execution inefficiency rather than a hardware resource limit—such as serial dependencies, poor instruction mix, or insufficient parallelism—and the remedy lies in how the kernel is written or selected.
2. **Near the bandwidth ceiling.** The kernel is saturating memory bandwidth; only transferring fewer bytes (through operator fusion, tiling, or a narrower numeric representation) can speed it up.
3. **Near the compute ceiling.** The kernel is saturating the arithmetic units; only performing fewer operations or switching to a faster supported numeric format or accelerator can reduce its duration.

Whereas the roofline narrows the set of candidate levers for an individual kernel, the step trace determines whether that kernel occupies enough of the critical path to warrant optimization.

A normalization kernel illustrates this three-step placement. Suppose hardware counters for a LayerNorm kernel on an H100 report 5 TFLOP/s of achieved compute and 2.8 TB/s of achieved HBM bandwidth. Its measured arithmetic intensity is 5 TFLOP/s ÷ 2.8 TB/s ≈ 1.8 FLOP/byte—more than two orders of magnitude below the H100's FP16 ridge point of about 295 FLOP/byte. The kernel therefore sits squarely on the memory-bandwidth slope, and its achieved bandwidth is already 83.6 percent of the 3.35 TB/s hardware peak. Faster arithmetic units cannot speed it up at all, and tuning its memory access pattern can recover at most the remaining sixth of headroom to the ceiling. Any substantial speedup must come from reducing D_{vol}: either by fusing the normalization into the preceding operator so the intermediate tensor never round-trips through HBM, or by adopting a narrower numeric format.

The kernels of our 70B training job divide along the same line. Large matrix multiplications in the attention projections and feed-forward blocks operate above the ridge point near the compute ceiling. By contrast, normalization, activation functions, softmax, residual additions, and the optimizer update sit far down the memory-bandwidth slope just like LayerNorm, executing a tiny fraction of the step's useful FLOPs while consuming a disproportionate share of local execution time. Comparing the time spent in matrix multiplications against the time spent in all other kernels therefore provides a rapid check on local program goodput losses. A step dominated by fast matrix multiplications loses its program goodput primarily to exposed communication and synchronization (as in our running example). A step spending substantial time in non-GEMM kernels, by contrast, has a local memory-traffic loss that operator fusion can recover.

The serving replica's decode step sits even farther down the bandwidth slope. At batch size one, each of its eight accelerators must stream its entire weight shard from HBM to emit a single token. That minimum streaming time, about 5.2 ms per token (section 2.3.1.3), serves as the named ideal for the replica's program goodput. Even at that ideal floor, the H100 tensor cores run at only 0.3 percent of their FP16 peak, rendering MFU uninformative. Instead, the replica's program goodput is measured by MBU: the ratio of the weight-streaming floor time to the measured step time (section 9.2.2).

Because the weight-streaming floor counts only model parameters, a measured decode step always takes longer. Reading the key-value cache, executing the per-layer tensor-parallel AllReduce, and dispatching dozens of small kernels per layer all lengthen the step and lower MBU, appearing as program goodput losses of the replica. How large a batch the replica can accumulate to climb the roofline slope toward the ridge point is governed by its latency service-level objective in chapter 10.

9.6.4 Common bottleneck patterns

An accelerator that idles *between* steps waiting for its next input batch shows the same drop in utilization as an accelerator that idles *within* a step waiting for a collective to finish. Yet the first is losing runtime goodput and answers to the data pipeline, whereas the second is losing program goodput and answers to communication overlap. Recognizing recurring trace signatures provides fast diagnostic hypotheses (table 9.7), and identifying which goodput factor a signature lowers is just as important as identifying its mechanism, since not every slow-step symptom is a program loss.

Table 9.7: Common Bottleneck Trace Patterns: Each trace signature maps to a likely cause, the goodput factor and term the cause lengthens, a candidate lever, and the trade-off to check before applying it. Levers marked with a section reference are developed there.

Trace pattern	Likely cause	Factor and term	Candidate lever	Trade-off to check
Small gaps between many short kernels	Host dispatch of each kernel dominates	Program; dispatch part of L_{lat}	Graph replay, fusion, or compilation (section 9.7.3)	Compilation time, charged to runtime goodput; variants per input shape
Large memory spikes	Unfused operators materialize intermediates in HBM	Program; D_{vol}/BW and HBM capacity	Fusion, or tiled attention that never writes its score matrix (section 9.7.3)	Numerical validation and kernel maintenance
Low FLOP/s with high achieved bandwidth	Memory-bound operations dominate the step	Program; D_{vol}/BW	Fewer bytes per step, through fusion, a smaller representation, or a larger per-device batch where memory and convergence allow	Accuracy, cache pressure, and quality guardrails
Much of the step in non-GEMM kernels	Normalization, activation, softmax, and optimizer kernels sit on the bandwidth slope	Program; D_{vol}/BW and dispatch	Fuse non-GEMM work into neighboring kernels (section 9.7.3)	Fusion legality and kernel maintenance
Collective bars during device idle gaps	Communication serialized rather than overlapped	Program; $T_{\text{comm}}(N) - T_{\text{overlap}}$	Overlap collectives with independent work (section 9.7.2); compiler-chosen partitioning	Bucket size, contention for device resources, schedule complexity

Trace pattern	Likely cause	Factor and term	Candidate lever	Trade-off to check
All ranks but one wait at a barrier	Uneven work or a degrading device	Program; $T_{\text{sync}}(N)$	Even partitioning; expert capacity limits for MoE (section 6.5.5); node replacement by the reliability layer	Restart cost, charged to runtime goodput, against the recovered step time
Stages idle at the start and end of each step	Pipeline bubbles	Program; $T_{\text{sync}}(N)$	More microbatches or an interleaved schedule (section 6.5.2)	Activation memory and per-microbatch efficiency
Repeated large allocations	Buffers recreated rather than reused	Program; allocator time in L_{lat}	Memory planning by a compiler (section 9.7.3) or preallocated buffer pools	Allocator behavior; recomputation cost if activations are dropped instead
Periodic activity drops between steps	Input preparation not pipelined with the step	Runtime; input wait, not a program loss	More data workers, preprocessing services, or pretokenized data (chapter 4)	Storage cost and data freshness

The final row of table 9.7 is the most frequently misdiagnosed in practice; checking the per-step input-wait field in the job timeline distinguishes an input stall from an exposed collective before any profiler trace is opened. Once the hierarchical drill-down concludes, it hands off a specific Fleet Law term and its share of the step—exposed cross-node communication for our training job, whose share the constructed example assumes, and the gap above the weight-streaming floor for our serving replica—which section 9.7 now recovers through targeted engineering levers.

9.7 Recovering Program Goodput

Tuning a single-device kernel when a step is bound by cross-node gradient reductions, or overlapping collectives when the step is bound by host kernel launches, leaves the critical path untouched. Recovering program goodput requires changing how a retained step executes, and an optimization helps only to the extent that it shortens an exposed term of the step. Three families of levers target distinct Fleet Law terms:

1. **Compiler-derived partitioning** determines which collectives a step executes and across which physical links they travel ($T_{\text{comm}}(N)$).
2. **Communication overlap** hides a portion of those collectives behind independent computation, enlarging the overlap credit (T_{overlap}).
3. **Kernel-level optimizations** (fusion, tiling, graph replay, and compilation) compress single-device execution (T_{compute}/N).

Because the speedup from any one lever depends on what the others leave exposed on the critical path, engineers apply these levers one at a time and re-measure the step after each change.

9.7.1 Compiler-derived partitioning

A model whose tensor-parallel modules place AllReduce calls directly inside the forward and backward methods of each layer (chapter 6) cannot change how those layers are split, or move to a node topology that calls for a different split, without modifying the model code itself. A data-parallel wrapper adapts to a new device count on its own, but hand-partitioned layers do not. Because how a large training job is partitioned across devices determines most of the communication its step performs, tying that layout to hand-written layer code makes every layout change a code change. Compiler-derived partitioning, pioneered in the Accelerated Linear Algebra (XLA) compiler's GSPMD partitioner (Xu et al. 2021), shifts the derivation of tensor layouts and the insertion of collective operations from the model author to the compiler.

Definition 9.5: Compiler-derived partitioning

Compiler-derived partitioning is the division of one program across a mesh of devices in which the programmer annotates how a few tensors are split along named mesh axes and

the compiler derives the layout of every other tensor and inserts the collectives those layouts require.

1. **Significance**: The inserted collectives are the step's communication term, so the annotations and the mapping of mesh axes onto the fleet decide how much of each retained step a job spends communicating. Because the layout is stated apart from the model code, a different layout is a different set of annotations rather than a rewritten program.
2. **Distinction**: Unlike programmer-specified parallelism, in which the engineer writes or selects each collective through a data-parallel wrapper or a hand-partitioned layer, compiler-derived partitioning leaves the engineer only the key annotations and derives the rest. The partitioning strategies themselves, data, tensor, and pipeline parallelism, are unchanged; what changes is who writes their communication.
3. **Common pitfall**: A frequent mistake is to read the step's communication from the annotations. Where derived layouts disagree, the compiler inserts resharding collectives that no line of the program names, so the step trace, not the source, is the record of what the step communicates.

Under this model, the engineer specifies only two inputs. The first is a logical device mesh whose axes carry symbolic names. For our running example, which uses tensor parallelism within each node and data parallelism across the 16 nodes, the mesh has a `model` axis of 8 devices and a `data` axis spanning the sixteen nodes. The second input is a small set of sharding annotations on key tensors, declaring which tensor dimension is partitioned along which mesh axis. The input batch dimension is split along the `data` axis. Within each feed-forward block, the gate and up projection weights (the two input matrices of the gated feed-forward layer, whose outputs the gated activation combines) are split column-wise along the `model` axis while the output projection weight is split row-wise. The attention projections are split by head along the `model` axis. Every other intermediate tensor in the program remains unannotated.

From those few seeds, the compiler derives the rest of the parallel program through sharding propagation, walking the computation graph to infer the output layout of each operator from the layouts of its inputs. Multiplying an input activation by a column-split weight matrix yields an activation split along its feature dimension; the subsequent element-wise gated activation preserves that layout without communication. Next, multiplying that feature-split activation by a row-split output weight produces a partial sum on each device, prompting the compiler to insert an AllReduce along the `model` axis to complete the projection. The backward pass mirrors this derivation. Each device's weight gradients cover only its local shard along the `model` axis but represent partial sums across the `data` axis, so the compiler automatically inserts a cross-node gradient reduction along the `data` axis (or a ReduceScatter paired with a forward AllGather when optimizer state is sharded). The resulting communication schedule matches the hand-partitioned Megatron-style layer of section 6.5.3, yet it is synthesized automatically from single-device model code that contains no explicit collective calls (figure 9.8). Wherever two incompatible tensor layouts meet in the graph, propagation resolves the mismatch by inserting an AllGather, ReduceScatter, AllToAll, or point-to-point permutation.

Every accelerator executes the same compiled binary, using its own mesh coordinates at runtime to index its local weight shard and select its collective peers—the single program, multiple data (SPMD) execution model from which GSPMD takes its name (Xu et al. 2021). Because a single compilation serves all ranks, compilation time scales with the size of the model graph rather than the number of devices in the fleet. The collectives that propagation inserted are part of that compiled binary, so a stored artifact carries its communication schedule to every later job or replica that runs the same model, mesh, and annotations. Changing the mesh dimensions or sharding annotations still triggers a fresh compilation, whose runtime goodput cost is weighed against its program goodput gain in section 9.9.3.

↩ Which strategies a training job combines, and the feasibility of each combination, are established in section 6.6.

Because the device mesh is a logical abstraction, the cluster scheduler's physical placement determines which hardware links carry each mesh axis. When an eight-device `model` axis maps onto a single node, its per-layer AllReduce collectives traverse high-bandwidth NVLink; if that same axis is scattered across nodes, those identical collectives traverse the inter-node fabric at a

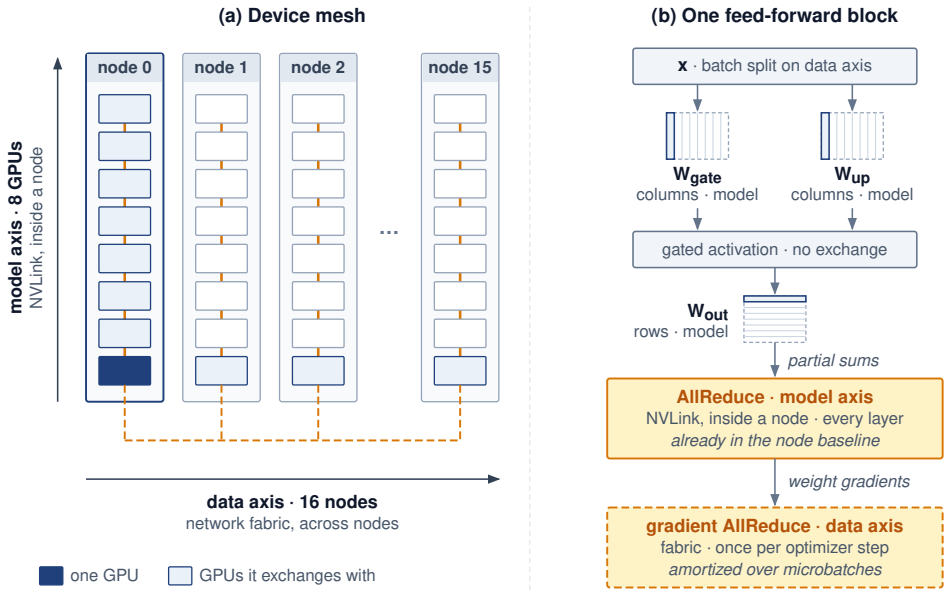

Figure 9.8: From Annotations to Collectives: The running example's mesh places the model axis of 8 GPUs inside each node and the data axis across the 16 nodes. In one feed-forward block, the column-split gate and up projections need no exchange, and the row-split output projection leaves partial sums that propagation completes with a model-axis AllReduce in every layer, over NVLink and already inside the node baseline. The weight gradients are reduced along the data axis across the fabric (a ReduceScatter and an AllGather when the optimizer state is sharded) once per optimizer step, after gradient accumulation over many microbatch steps. The reduction is too long to fit inside one illustrative microbatch step and is amortized across the accumulated microbatches, while the time each microbatch step spends beyond the node baseline is assumed to be per-microbatch synchronization and collective waits. Schematic; representative nodes only.

fraction of the bandwidth. Three distinct mechanisms therefore govern a step's communication cost: sharding annotations determine *which* collectives are synthesized, physical mesh placement determines *how fast* each collective runs, and stream scheduling determines *how much* of that duration remains exposed. In Fleet Law terms, compiler-derived partitioning sets $T_{\text{comm}}(N)$ and shapes the opportunity for T_{overlap}, while leaving single-device execution (T_{compute}/N) to the kernel levers of section 9.7.3.

Consequently, modifying a sharding annotation alters the step's communication pattern without changing a single line of mathematical model code, and its effect must be verified under the measurement contract of section 9.3.1. The step trace provides the necessary audit: counting the collectives in one step by operation type and mesh axis, and checking that count against the intended layout, confirms whether propagation behaved as expected. An unexpected AllGather appearing in every layer is the classic symptom of an annotation conflict—typically caused by an intermediate reshape or a single misannotated tensor that forces the compiler to reshard back and forth. Because a compiler's internal cost model can also pick a suboptimal propagation path, any compiler-derived layout should be treated as a hypothesis about $T_{\text{comm}}(N)$ until confirmed by the step trace.

At the fleet level, decoupling layout from model code allows the partitioning to adapt seamlessly to the allocation. When the fleet orchestrator resizes an elastic job (section 8.7.3) onto a different node topology, the `model` axis can change along with the `data` axis, and the compiler regenerates the executable from the existing annotations. Exploring layout trade-offs—such as sharding optimizer states or shifting communication onto intra-node links—likewise reduces to editing annotations rather than rewriting layers. At the same time, our running example illustrates the boundary of what partitioning alone can achieve. Its mesh already confines every `model`-axis collective within a single node, so the extra time its 128-GPU step spends beyond the single-node baseline comes from traffic along the cross-node `data` axis. Whether that traffic can be hidden depends on how

its transfer time compares with the backward compute window, and for the largest exchange, the gradient reduction, the answer decides when the reduction runs at all.

9.7.2 Tensor-parallel collectives under contention

Once the partitioning layout fixes which collectives a step must execute, the exposed communication time is whatever fraction of those collectives the schedule fails to hide behind independent computation. Our running example executes two classes of collectives that differ in their data dependencies. Cross-node gradient reductions along the `data` axis have loose dependencies. As soon as the backward pass finishes computing the gradients for one bucket of layers, that bucket can be reduced across nodes while backward computation continues on earlier layers. This is the standard bucketed overlap implemented by `DistributedDataParallel`. The analytical limits on hiding such loosely coupled exchanges are established in section 5.10.4, and for our 128-GPU job the full gradient reduction lies far outside them.

Each GPU holds the 16-bit gradients of its eighth of the 70B parameters, 17.5 GB, and a ring reduction across the 16 nodes sends 1.875 times that shard through the GPU's own network interface. Even at the full 50 GB/s of an NDR port, the transfer takes at least 656 ms, 2.7× the entire 245 ms microbatch step, so no bucket schedule can hide it inside one step. Sharding the optimizer state does not rescue a single step, because the ReduceScatter that must overlap with backpropagation still needs about 328 ms, longer than the whole step.

The job therefore accumulates gradients over many microbatch steps and reduces them once per optimizer step. Holding a fully exposed reduction to 5 percent of the optimizer step requires accumulated computation 19 times as long as the reduction, about 12.5 s, which takes 51 microbatch steps of 245 ms, or about 1.8 million tokens per update. The archetype's 175B reference job makes the same trade at larger scale, accumulating 96 microbatches per replica and spending about 2 percent of its optimizer step on the data-parallel reduction. At 51 microbatches, the reduction adds about 13 ms per microbatch on average, time that lies outside the microbatch step behind the chapter's MFU figures and lowers program goodput over a whole optimizer step from 0.48 to about 0.456. The gradient reduction therefore cannot be what fills the 65 ms that each illustrative microbatch step spends beyond the node baseline. The constructed example assumes that this time is per-microbatch synchronization across the 16 nodes and the collective waits it causes, which recur on every microbatch and cannot hide; the example does not size them as a particular operation.

By contrast, intra-node collectives along the `model` axis have strict serial dependencies: each AllReduce completes an activation tensor that the very next operator in the same transformer layer consumes immediately, preventing subsequent operations in that layer from starting until the collective finishes. Because these `model`-axis exchanges never leave the node, their duration is already embedded inside the single-node baseline.

A bandwidth-floor estimate (excluding launch latency and activation recomputation) bounds how much of the single-node step those intra-node collectives can occupy. Each exchange point reduces roughly 36 MB of activations across the 8 GPUs of the node, and the 320 exchange points in a full forward-and-backward step require about 45 ms of NVLink transfer time. Even if none of that transfer were hidden, it would represent about 25 percent of the 180 ms single-node step—well within the 35 percent gap between the node's 65 percent MFU and the arithmetic ideal. Consequently, `model`-axis collectives represent a local, node-level program loss that does not grow with fleet size so long as the `model` axis remains confined to a single node. The cross-node `data`-axis traffic carries the scaling penalty, and a single-rank trace easily separates the two because they execute on distinct process groups.

Hiding a collective whose output is needed by the next operation requires finding independent computation to execute concurrently. One approach interleaves work from a second microbatch that has no dependency on the collective in flight. The alternative is fine-grained compute-communication decomposition. It splits both the matrix multiplication and the collective that consumes its output into smaller tiles, so that one tile's collective transfer runs concurrently with the next tile's matrix multiplication. A partitioning compiler can synthesize this transformation automatically (Wang et al. 2022). Neither strategy shortens the communication time $T_{\text{comm}}(N)$ itself; both work purely by enlarging the overlap credit T_{overlap}.

However, overlapping communication and computation on the same accelerator is never contention-free. On NVIDIA GPUs, concurrent execution is orchestrated through CUDA streams—ordered queues of kernels, memory copies, and collectives that execute sequentially within a stream but may run concurrently across distinct streams. A distributed training step typically assigns matrix multiplications and element-wise ops to a compute stream while dispatching NCCL collectives to a separate communication stream (listing 9.1).

Listing 9.1: Computation and a Collective on Separate Streams: Both streams first wait for the stream that produced their inputs, and each tensor is recorded on the stream that uses it so that its storage is not reused while that stream still needs it. The matrix multiplication and the AllReduce may then run concurrently. Concurrency on the timeline does not by itself establish a shorter step.

```
compute_stream = torch.cuda.Stream()
comm_stream = torch.cuda.Stream()
producer_stream = torch.cuda.current_stream()
# Inputs are ready on producer_stream; gradients do not alias A or B.
compute_stream.wait_stream(producer_stream)
comm_stream.wait_stream(producer_stream)
A.record_stream(compute_stream)
B.record_stream(compute_stream)
gradients.record_stream(comm_stream)

with torch.cuda.stream(compute_stream):
    output = torch.matmul(A, B)  # Non-blocking compute

with torch.cuda.stream(comm_stream):
    work = dist.all_reduce(gradients, async_op=True)
    work.wait()  # Establish completion on the communication stream

torch.cuda.synchronize()  # Wait for both to complete
```

Placing work on separate streams enables concurrency, but it does not partition the underlying hardware resources. Because collective libraries such as NCCL launch their own GPU kernels to copy, pack, and reduce buffers across NVLink and network interfaces, an overlapped collective competes directly with concurrent matrix multiplications for streaming multiprocessors (SMs), L2 cache, and HBM bandwidth (Pati et al. 2024). In a kernel timeline, successful stream overlap shifts the collective bar underneath the compute kernels, but resource contention stretches the duration of those compute kernels, causing the step to finish later than an idealized contention-free model would predict (figure 9.9). Overlapping a collective improves program goodput only if the communication time removed from the critical path exceeds the contention delay added to the concurrent computation—a net balance that only an end-to-end step measurement can verify.

When a program manages multi-stream overlap manually rather than relying on a tested wrapper, it must also enforce two correctness invariants. Data readiness uses `wait_stream` so that no consumer stream reads a tensor before its producer stream has finished writing it. Storage lifetime uses `record_stream` so that the caching allocator does not recycle a tensor's memory buffer while a side stream is still reading it. Because host kernel launches are asynchronous, a tensor's Python reference can go out of scope on the CPU before the side stream finishes using the buffer on the GPU, and the allocator tracks buffer liveness only on the allocating stream unless `record_stream` registers the side stream as well. Both invariants must be verified before crediting a shorter step time to communication overlap rather than to a data race that silently corrupted or skipped work.

Our 8-GPU serving replica, partitioned purely along the `model` axis within a single node, illustrates the opposite regime. At batch size one, each per-layer AllReduce transfers only about 32.8 KB across the 8 GPUs—orders of magnitude below the message size required to exit the latency-bound regime of section 5.5. At this scale, fixed collective launch and protocol overhead (roughly 70× the wire transfer time itself) dominates the exchange, so a faster link would shorten only the small transfer portion. Stream overlap cannot hide the exchange either, because the immediately following layer operator awaits the reduced activation. Shortening this term requires either reducing the number of

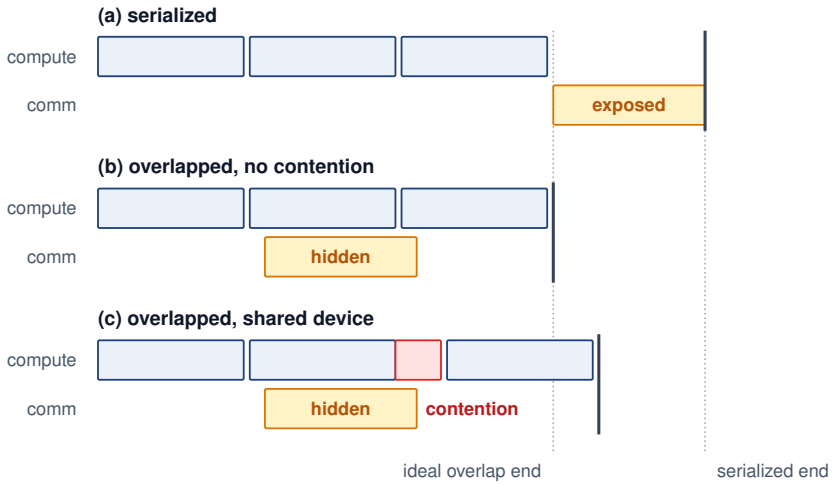

Figure 9.9: Overlap on a Shared Device: One step of three kernels and one collective on a compute stream and a communication stream. Serialized, the collective runs after the kernels and is exposed. Overlapped on a device with resources to spare, it hides completely behind the kernels. Overlapped on a real device, the collective's own kernels take streaming multiprocessors and memory bandwidth from the kernel beside it, which runs longer, so the step ends between the two. Only a measured step shows where between them it ends. Durations are schematic, not measured.

exchanges through the partitioning layout or cutting per-launch overhead via CUDA Graph replay. How larger decode batches alter this balance is analyzed in section 10.10.2.

9.7.3 Kernel-level levers

Even when inter-device communication is well hidden, an accelerator can waste much of its local execution time writing short-lived intermediate tensors out to HBM and reading them immediately back, materializing quadratic attention score matrices, or sitting idle between thousands of small host-dispatched kernels. Four single-device optimization techniques—operator fusion, tiling, graph replay, and graph compilation—recover these local losses without altering the model's mathematical semantics. In the Fleet Law, all four techniques compress only the local execution term T_{compute}/N; because none of them directly reduces $T_{\text{comm}}(N)$ or $T_{\text{sync}}(N)$, their end-to-end fleet impact diminishes whenever exposed communication or barrier waits dominate the step.

↪ Fusion, tiling, graph replay, and compilation are collected for reference in section A.5.

Operator fusion merges a chain of consecutive operations into a single GPU kernel so that intermediate activations remain in registers or shared memory (SRAM) instead of round-tripping through HBM, while replacing multiple host kernel launches with one. It therefore reduces both the memory traffic volume D_{vol} and the host-dispatch portion of L_{lat}. For a linear layer followed by an activation and normalization at the hidden dimension of Llama 2 70B and a batch of 2,048 tokens, eliminating the HBM write-and-read of the two intermediates cuts total modeled HBM traffic by a factor of about 1.7×. The hardware limit on fusion is register and SRAM pressure: keeping too many live intermediates per thread reduces SM occupancy and can eventually slow the kernel down. In a profile, successful fusion manifests as fewer kernel launches, fewer measured HBM bytes transferred, and a shorter step duration (section A.5.1).

The boundary comparison in figure 9.10 shows which intermediate transfers fusion can remove.

A profile must confirm that the intermediate transfers disappear. Fewer launches alone establish a dispatch change, not a reduction in HBM traffic.

Tiling partitions a computation into blocks sized to fit in on-chip SRAM so that each operand fetched from HBM is reused across many arithmetic operations before eviction. Applied to self-attention (as in FlashAttention (Dao et al. 2022)), tiling evaluates attention block by block using online softmax rescaling, avoiding ever materializing the quadratic sequence-by-sequence score matrix in HBM. For a single attention head of our 70B model at a sequence length of 8,192 tokens, an idealized traffic bound shows that eliminating the score and probability matrix round-trips

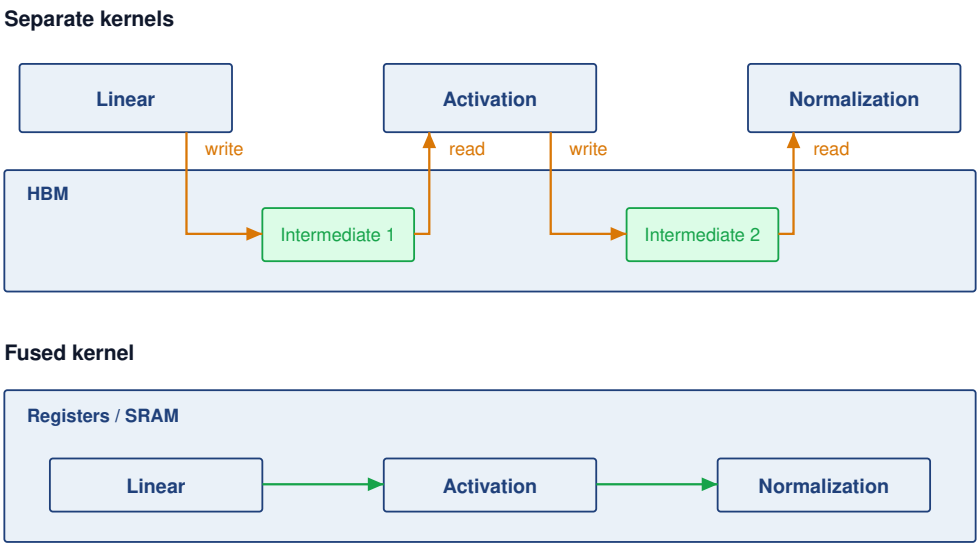

Figure 9.10: Fusion Across the HBM Boundary: Separate kernels write each intermediate tensor to HBM and read it back for the next operation. A fused kernel keeps these intermediates in registers or SRAM as it executes the same chain. The diagram isolates intermediate traffic; common input, weight, and final-output transfers remain outside the drawing. The fused schedule assumes compatible operations and enough on-chip storage for the live working set, including normalization state.

reduces HBM traffic by roughly 65×. Beyond cutting D_{vol}, tiling frees substantial HBM capacity, which training jobs reinvest in longer context lengths or larger microbatches and serving replicas reinvest in larger key-value caches. The trade-off is extra arithmetic for online rescaling and, during backpropagation, recomputing attention scores on the fly rather than storing them. Success is verified by the disappearance of quadratic score-matrix allocations and HBM traffic in the memory trace (section A.5.2).

The working set in figure 9.11 stays local while attention consumes successive key/value tiles.

The memory saving comes from the allocation that never crosses the HBM boundary, not from omitting attention pairs. Tile size must leave room for the retained state as well as the operands.

Graph replay captures a fixed sequence of GPU kernel launches once and replays the entire subgraph with a single host API call, compressing the host-dispatch portion of L_{lat} while leaving kernel execution itself unchanged. In eager execution, dispatching the kernels of a single forward pass of our 70B model consumes roughly 12–24 ms of host CPU overhead—several times longer than the 5.2 ms weight-streaming floor of the serving replica's decode step (section 9.6.3). Eager small-batch decode therefore becomes host-dispatch bound long before it saturates HBM bandwidth. Distributed training encounters the same regime under strong scaling: when a job expands across more accelerators at a fixed global batch size, per-device kernel durations shrink while the number of kernel launches per step stays constant. The constraint of graph replay is static structure: tensor shapes, memory addresses, and control flow must remain fixed across replays, requiring serving engines with dynamic batch sizes to capture and cache a separate graph for each padded shape bucket. Success is verified when inter-kernel launch gaps vanish from the kernel timeline (section A.5.3).

Graph compilation traces a region of the model into an intermediate computation graph, rewrites the graph, fuses compatible operators, plans static buffer reuse, and generates or autotunes specialized kernels for each fused subgraph—automating fusion, memory planning, and kernel selection in a single pass. Compilation reduces D_{vol} and host dispatch L_{lat}, and when autotuning selects superior tile shapes it increases hardware compute efficiency η_{hw} as well. Because compilation runs before the first retained step (and, without a persistent compilation cache, again after each restart, preemption, or elastic resize to a new mesh), its build duration is a runtime goodput cost that must be amortized across subsequent execution, a cross-factor trade-off quantified in section 9.9.3. A recompilation triggered by new input shapes during steady-state execution instead slows retained

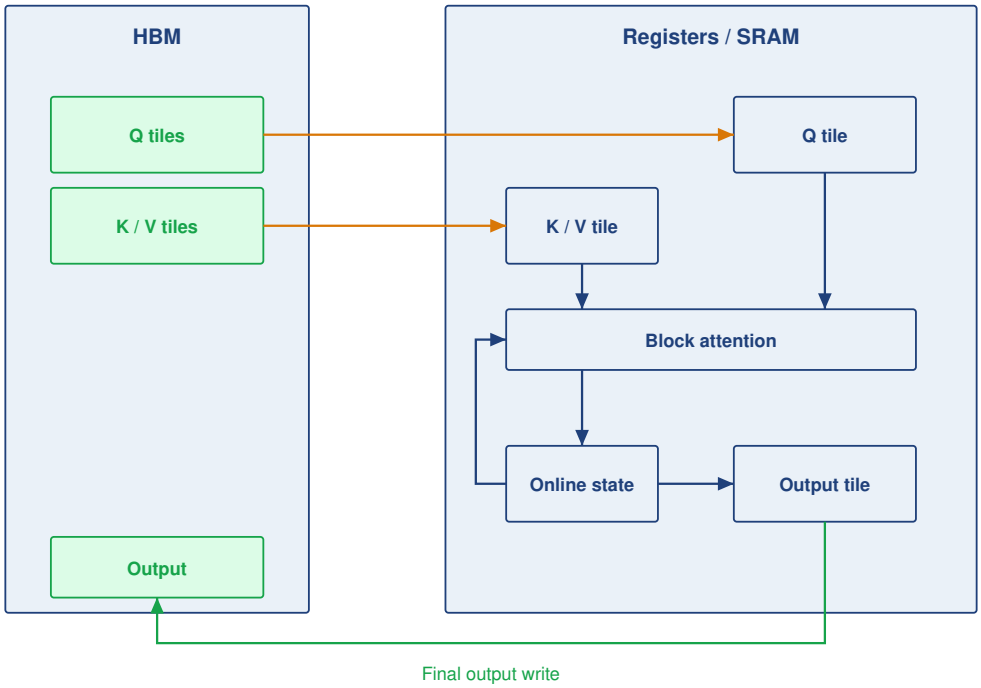

Figure 9.11: Tiled Attention and Online State: For one query tile, the kernel reads relevant key/value tiles from HBM and updates on-chip state after each block. This state contains running row maxima, normalization sums, and output accumulators; the feedback path carries it into the next block. After the last required block, the kernel normalizes and writes the output tile. Score tiles remain temporary on-chip values rather than a sequence-by-sequence HBM allocation. Another query tile can require reading the same key/value tiles again; the diagram does not imply one global read.

steps and lowers program goodput. Success is verified by fewer, larger kernels in the timeline and a shorter steady-state step measured after compilation completes (section A.5.4).

Numeric precision also acts on local execution, although unlike the four techniques above it changes the model's arithmetic. This book divides its mechanics between two domain chapters: reduced-precision training (such as FP8), which raises the compute ceiling R_{peak} and halves activation and weight traffic under the numerical stability contract of section 6.4.8, and post-training weight quantization for inference, which reduces weight-streaming bytes D_{vol} in section 10.9.

These single-device levers pay off very differently across our two running workloads. In the 128-GPU training job, large matrix multiplications already run near the H100 compute ceiling, so kernel-level levers can compress only the non-GEMM kernels on the bandwidth slope while leaving the job's primary program loss, exposed cross-node communication, untouched. In the 8-GPU serving replica, by contrast, the program loss is overwhelmingly local (weight-streaming traffic and host kernel dispatch), allowing operator fusion, weight quantization, and CUDA Graph replay to close part of the gap to the bandwidth floor, as the replica case study in section 10.15.1 shows.

9.7.4 Composing program fixes

Applying several kernel and communication optimizations simultaneously—or adding their isolated microbenchmark speedups on a spreadsheet—almost always overpredicts the fleet gain, because speeding up backward compute kernels shrinks the execution window that previously hid gradient communication. Effective performance engineering therefore follows a disciplined, sequential loop that starts from the whole job rather than an individual kernel:

1. Measure the program goodput of retained steps under the measurement contract (section 9.3.1), and use factor localization (section 9.3.2) to confirm that the bottleneck lies in program goodput rather than scheduling or runtime goodput.
2. Descend the profiling hierarchy (section 9.6.3) to identify which Fleet Law term holds the exposed time on the critical path, and select the lowest-cost candidate lever targeting that term (table 9.7).
3. Apply a single lever, re-measure the end-to-end step under the exact same contract, and restart the drill-down from the new trace—because compressing one exposed term routinely exposes another.

For our 128-GPU training example, the constructed drill-down identified exposed cross-node synchronization and collective waits along the `data` axis as the primary bottleneck. Where an exchange's transfer time fits within the backward compute window, the first experiment targets communication overlap, either by tuning gradient bucket sizes and launch order or by sharding the optimizer state so that ReduceScatter overlaps with backpropagation and AllGather prefetches parameters ahead of the forward pass. The full gradient reduction does not fit. Even with the optimizer state sharded, its ReduceScatter alone needs about 328 ms, longer than the whole 245 ms microbatch step, so overlap cannot hide it within a step. The lever for that reduction is gradient accumulation. Running it once per optimizer step, after about 51 microbatches (section 9.7.2), holds even a fully exposed reduction to 5 percent of the update, and accumulating more microbatches shrinks that share further at the cost of a larger global batch. Sharding then reduces what must overlap with the final backward pass rather than making the reduction fit. Only after re-profiling that change does the second experiment target the remaining node-level losses (intra-node `model`-axis collectives and memory-bound non-GEMM kernels) via fine-grained compute-communication decomposition and operator fusion.

Isolated speedups fail to add linearly because every optimization alters the step's critical path. A lever shortens a single Fleet Law term, and end-to-end step time decreases only by the portion of that term that was exposed. Once the first lever takes effect, every term's share of the step shifts, invalidating isolated speedup estimates for subsequent levers. In practice, pairs of optimizations interact in three ways:

- **Reinforcing combinations.** Operator fusion and tiled attention free HBM activation memory; reinvesting that freed memory in a larger microbatch lengthens the compute window per step, which in turn hides a larger fraction of collective communication.
- **Redundant combinations.** Enabling a graph compiler that automatically fuses element-wise operator chains leaves nothing for a hand-written fused kernel on the same chain to recover.
- **Conflicting combinations.** Overlapping NCCL communication kernels steals SMs and HBM bandwidth from concurrent matrix multiplications, lengthening local compute. Conversely, speeding up backward compute kernels via FP8 or fusion shrinks the compute window that hid cross-node gradient reductions, exposing communication that was previously hidden for free.

Because of these interactions, an optimization validated on a single-node baseline must always be re-evaluated on the full fleet configuration. Only an end-to-end measurement under table 9.2 can establish whether a combination of levers is reinforcing, redundant, or conflicting—and whether a rise in program goodput came at the expense of scheduling or runtime goodput (section 9.9).

Checkpoint 9.3: Composing program fixes

A program fix recovers goodput only through the exposed term it shortens, and only a new measurement shows what it added.

Choosing the lever

- ☐ **Decode at low MBU**: The serving replica's decode step runs well above its bandwidth floor. Can you name the trace evidence that would separate a dispatch-bound step,

a step that streams more bytes than the weights require, and a step that exposes its tensor-parallel exchanges, and the lever each points to?

- ☐ **FP8 on a communication-bound step**: Can you explain why faster lower-precision arithmetic alone may leave the running example's fleet step nearly unchanged, and identify what would have to be true of its exposed term for lower precision to help?
- ☐ **Annotation change**: After an annotation change, which collectives should a step trace show, on which mesh axis, and what does an unexpected AllGather in every layer indicate?

Composing and measuring

- ☐ **Non-additive speedups**: Two fixes each shorten the node step when measured alone. Can you describe how applying both could yield less than the sum of their savings, and one way it could yield more?

9.8 The Whole Fleet

The Llama 2 70B job keeps 0.90 of its held time as retained steps, runs those steps at 0.48 of the ideal rate, and draws its nodes from a pool that holds 0.65 of its capacity in complete allocations. Each figure is accurate on its own boundary, yet none states alone how much of the fleet's capacity becomes training progress or which engineering lever yields the largest return. Those answers emerge only from the product of equation 9.1, which also converts a fraction of fleet capacity into the unit cost of useful work.

9.8.1 Decomposing the running example

Evaluated solely by its runtime and program goodput inside its own allocation, the Llama 2 70B job appears nearly half efficient, yet the shared pool that supplies its nodes loses roughly another third of its capacity before any job holds a complete gang. Composing all three factors for one job on one fleet reveals where the entire efficiency gap goes across the cluster.

Napkin Math 9.5: Where the running example's fleet capacity went

Problem: The Llama 2 70B job holds 128 GPUs drawn from a shared pool. What share of the accelerator-hours the job holds becomes progress at the ideal rate, what share of the capacity the pool offers becomes progress, and at which boundary is the rest lost?

Variables:

- Scheduling goodput $\eta_{\text{scheduling}} = 0.65$, the pool's value from notebook 9.2.
- Runtime goodput $\eta_{\text{runtime}} = 0.90$, composed from five owners' terms in notebook 9.3.
- Program goodput $\eta_{\text{program}} = 0.48$, the per-microbatch fleet MFU of the retained steps from notebook 9.4; over a full optimizer step the amortized gradient reduction lowers it to about 0.456 (section 9.7.2).

Math:

- Job product, $\eta_{\text{runtime}}\eta_{\text{program}}$: $0.90 \times 0.48 \approx 0.43$.
- Pool product, MPG $= \eta_{\text{scheduling}}\eta_{\text{runtime}}\eta_{\text{program}}$: $0.65 \times 0.43 \approx 0.28$.
- Loss shares by equation 9.2: scheduling $1 - 0.65 = 0.35$; runtime $0.65 \times 0.10 \approx 0.06$; program $0.59 \times 0.52 \approx 0.31$, where 0.59 is $\eta_{\text{scheduling}}\eta_{\text{runtime}}$. With the 0.28 that became progress, the four parts add to one.
- Calibration: Llama 3 405B's product of about 0.37 on its dedicated cluster (notebook 9.1) becomes $0.65 \times 0.37 \approx 0.24$ on this pool.

Result: The job turns about 43 percent of the accelerator-hours it holds into ideal-rate progress, and the pool turns about 28 percent of its offered capacity into progress. Of that offered capacity,

35 percent is lost before any job holds it, 6 percent is held without producing a retained step, and 31 percent is spent on retained steps that run below the ideal rate.

Systems insight: The scheduling layer multiplies every job's product by the same 0.65, so even the best-documented large run on record, moved onto this pool, would deliver only about a quarter of the pool's capacity as progress. No engineering inside a job recovers a loss that occurs before the job holds its devices.

Both products are meaningful (table 9.8), and they answer different questions. The job's product ($\eta_{\text{runtime}} \times \eta_{\text{program}}$) measures what the team running the job can improve directly: runtime and program losses occur inside its allocation, and a higher job product means fewer held hours per retained step and an earlier result. The pool product (MPG) measures what the fleet's owner pays for: capacity that never forms a complete allocation is still powered, depreciated, and billed. Raising the pool's scheduling goodput does little for an individual job that already holds its allocation, because the capacity it recovers becomes complete allocations for other jobs waiting in the queue. Reporting only the job's product hides the pool's largest loss from the infrastructure team, whereas reporting only the pool's product charges the workload team for a scheduling loss it cannot control.

The pool product also rests on an assumption that the job product does not require. Because it multiplies the pool's scheduling goodput by this job's runtime and program goodput, it describes a pool in which every workload runs like the running example. The pool-level runtime losses identified in section 9.5.2—reserved blocks that no job uses and zombie allocations—lower the pool's effective runtime goodput further, making the pool figure an upper bound.

Table 9.8: Goodput Decomposition of the Running Example: The three factors of the Llama 2 70B job and the share of fleet capacity each boundary loses, computed by equation 9.2 in the order a job meets the losses. The scheduling layer comes from a constructed snapshot and the program layer from the example's constructed microbatch-scale step times, so both are illustrative. Program goodput is the per-microbatch value; the data-axis gradient reduction, amortized over the accumulated microbatches, lowers it to about 0.456 over a full optimizer step and MPG from about 0.282 to about 0.268. The Llama 3 column multiplies factors reported over different measurement windows and is an estimate (Dubey et al. 2024). The pool product assumes every job in the pool runs like the running example, so it is an upper bound. The serving replica's program goodput divides by a bandwidth ideal rather than peak arithmetic, so a fleet report carries it as its own row and never averages it with the training job's.

Boundary	Factor value	Share of fleet capacity lost here	Llama 3 405B (cited)	Lever owner
Scheduling goodput (illustrative)	0.65	35%	1 (dedicated)	chapter 8
Runtime goodput	0.90	6%	at least 0.90	chapter 7, chapter 4, chapter 8
Program goodput	0.48	31%	about 0.41	section 9.7
Job product ($\eta_{\text{runtime}} \times \eta_{\text{program}}$)	0.43		about 0.37	
Pool product (MPG)	0.28	28% useful	about 0.24 on the same pool	
Serving replica (separate row)	Runtime from 0.996 (released daily) to about 0.91 (one burst hour); program goodput as MBU against a floor of 5.2 ms per token	Not added to training	Not applicable	chapter 10

Read as a map of where the fleet's capacity went, table 9.8 shows that two boundaries carry almost all of the loss. The capacity lost before allocation and the capacity lost inside slow steps are each several times larger than the capacity lost to startup, input waits, checkpoints, failures, and preemptions combined (figure 9.12). Those runtime interruptions are nevertheless the losses that operators see most directly in job logs and restart alerts.

Comparing the running example against the Llama 3 405B column highlights how cluster scale shifts the balance between factors. The running example's program goodput sits above the Llama 3 reference by construction, reflecting the expectation that a job on 16 nodes exposes less communi-

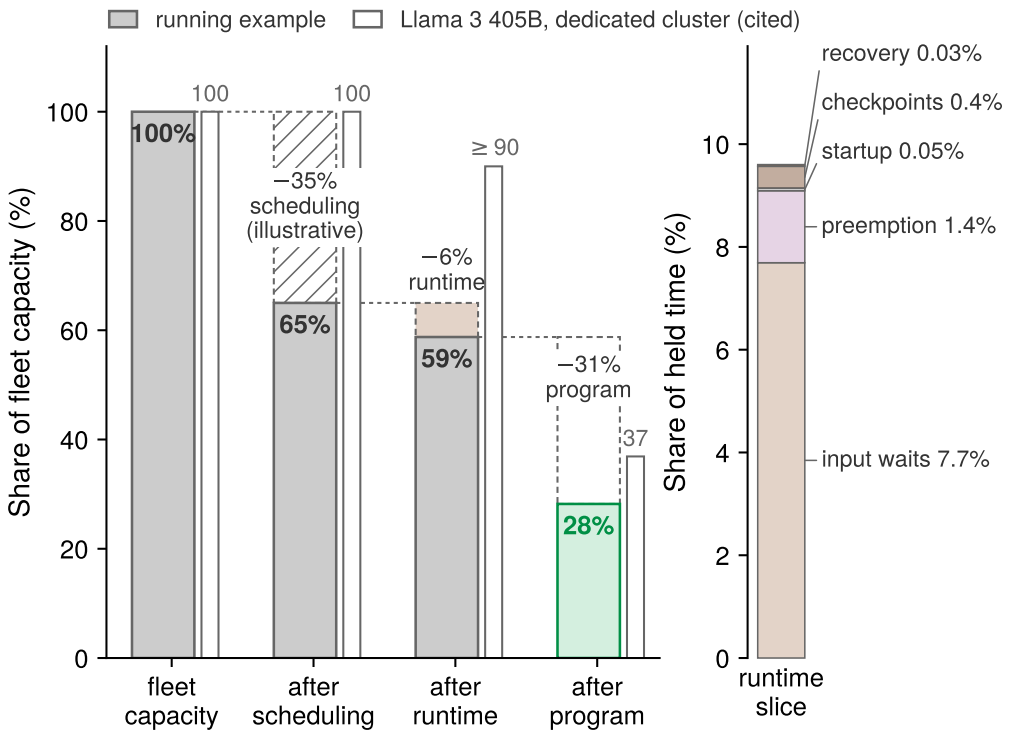

Figure 9.12: Where the Running Example's Fleet Capacity Goes: Each solid bar is the share of fleet capacity left after one boundary, and each dashed slice is the share lost at that boundary, so the three slices are the terms of equation 9.2: 35 percent lost to scheduling, 6 percent to runtime, and 31 percent to the program, leaving 28 percent as ideal-rate progress. The hatched scheduling slice comes from a constructed snapshot and is illustrative, and the program slice rests on the example's constructed microbatch-scale step times and uses the per-microbatch program goodput; amortizing the data-axis gradient reduction over the accumulated microbatches lowers it to about 0.456 over a full optimizer step and MPG from about 0.282 to about 0.268. The right panel opens the runtime slice into its five terms as shares of the time the job holds, which scheduling goodput scales down to the slice's share of fleet capacity; input waits and preemption carry almost all of it. The hollow bars place Llama 3 405B on its dedicated cluster beside the example, an estimate that multiplies factors reported over different measurement windows (Dubey et al. 2024).

cation per step than a run across 16,384 H100 GPUs. Its runtime goodput matches the reference, although a different mix of losses produces it, as section 9.5.3 analyzes.

Ranked by the share of the pool's capacity lost at each boundary, scheduling comes first, program goodput comes second, and runtime goodput comes a distant third. However, a boundary's loss share is not what an engineering lever recovers, because capacity recovered at an early boundary must still survive the later boundaries.

Instead, the product equation gives the marginal value of a lever directly. One percentage point of any factor is worth the product of the other two factors, because raising one factor by a small increment increases the overall product by that increment times the remaining two factors. For the running example, one point of scheduling goodput is worth 0.43 of a point of MPG, one point of runtime goodput is worth 0.31, and one point of program goodput is worth 0.59. A point of program goodput is worth the most because its multiplier is the product of the two highest factors, whereas a point of runtime goodput is worth the least because its multiplier is the product of the two lowest factors. Consequently, every factor's lever is worth more on a fleet whose other two factors are already high.

The gains a lever can achieve are also bounded by its own structural ceiling, and those ceilings reinforce the ranking. Removing every partial-gang hold and filling every freed device raises scheduling goodput to at most 0.80 (section 9.4.3). Eliminating input waits—the largest runtime term in section 9.5.2—raises runtime goodput to 0.98. Eliminating every loss the fleet adds beyond a

single node's baseline—mostly exposed communication (section 9.6.2)—raises program goodput to the single-node baseline of 0.65. Because each lever multiplies MPG by its factor's relative gain,[8] reaching those ceilings raises MPG by 23 percent (to 0.35) for the scheduling lever, by 9 percent (to 0.31) for the input pipeline, and by 35 percent (to 0.38) for the program lever. For this workload, the program lever yields the largest gain even though scheduling loses a larger share of the pool's total capacity, and the input pipeline yields the smallest gain of the three. Although these ceilings are optimistic—no communication overlap hides all exposed transfers and no packer fills every freed device—the priority order holds unless one team's practical reach falls far short of the others'. Still, this ranking is a property of the workload's factors, here constructed, rather than a universal law of the accounting; a fleet running a different model archetype will arrive at a different priority order.

[8] **Relative gains and boundary order**: Raising one factor multiplies MPG by the ratio of its new value to its old one, whatever the other two factors are, so the ranking of levers at their ceilings does not depend on the order in which the boundaries are drawn. The loss shares of equation 9.2 do depend on that order, because each later factor's loss is measured on what the earlier factors passed. A ranking by loss share can therefore differ from a ranking by relative gain, and only the second gives what a lever is worth.

Lighthouse 9.2: Archetype B (DLRM at Scale): Where a recommendation fleet loses goodput

A recommendation model stores embedding tables that far exceed accelerator memory in host DRAM or across remote shards (section 4.3.3). When the embedding rows required by a batch are fetched ahead of the training step, they belong to the input path; if that path cannot keep pace with the step rate, the accelerators stall waiting for their batch. The input-wait rule of section 9.2.3 charges that stall to runtime goodput, and the data stall ratio of section 4.4 sizes it. Inside a retained step, the gathered embedding vectors travel to their destination ranks through an AllToAll collective whose completion time is governed by the slowest destination under shard skew or fabric congestion (section 5.8). Whatever portion of that exchange is not hidden by dense computation lengthens the retained step and is charged to program goodput. Because a recommendation step is bounded more by memory and network bytes moved than by dense arithmetic, its program goodput must define its ideal denominator with the same care as a bandwidth-bound serving replica. The engineering levers shift accordingly: runtime goodput rises with prefetching and hot-row caching, whereas program goodput rises with shard load balancing and message coalescing (section 6.5.6). The marginal value of one point of any factor remains the product of the other two factors, but the lever ranking must be computed from the recommendation fleet's own measured numbers.

9.8.2 Cost per useful unit

An organization pays for fleet capacity by the accelerator-hour regardless of whether a complete allocation holds it, whether a retained step uses it, or whether that step runs at the ideal rate. Every factor of the goodput decomposition therefore divides directly into the cost of useful work. Specifically, the price of one accelerator-hour of ideal-rate progress equals the raw capacity price divided by MPG. With the capacity price per accelerator-hour $R_{\text{GPU/hr}}$ and the cost of one useful accelerator-hour C_{useful}, equation 9.6 states this relationship:

$$C_{\text{useful}} = \frac{R_{\text{GPU/hr}}}{\text{MPG}} = \frac{R_{\text{GPU/hr}}}{\eta_{\text{scheduling}} \times \eta_{\text{runtime}} \times \eta_{\text{program}}} \tag{9.6}$$

This relation turns each goodput factor into an inverse cost multiplier: a loss at any boundary raises the price of every unit of useful output in exact proportion. Using the internal fleet reference rate of \$2/GPU-hour from section 12.9 as the cost of one in-service accelerator-hour, dividing by scheduling goodput charges the scheduling loss once, consistent with the in-service denominator defined in section 9.4.1. Under this accounting, the running example's pool pays roughly \$7.09/GPU-hour for each accelerator-hour of ideal-rate progress, or 3.5× the raw capacity price. Because any unit of useful output—whether a training token or a served token—requires a fixed number of ideal-rate accelerator-hours, its unit cost carries the exact same multiplier. The spot-instance cost equation (equation 8.3) is the special case that divides the instance price by runtime goodput alone.

A chargeback policy that bills each workload team only for the accelerator-hours its jobs hold charges the running example 2.3× the raw rate per useful accelerator-hour and leaves the pool's scheduling loss unbilled. How an organization assigns that unallocated pool cost is a matter of cost-attribution policy (section 12.9.7.2).

Raising any goodput factor lowers the cost of every useful unit in proportion. Lifting the running example's program goodput to its single-node ceiling would reduce the cost of useful work by 26 percent without purchasing a single additional accelerator. However, the cost multiplier alone does not dictate which improvement to pursue, because levers that raise one factor frequently lower another; section 9.9 evaluates those cross-boundary trade-offs using the full goodput product.

9.9 Trading Between Factors

Waiting for a compact node placement, running a longer autotuning compilation before the first step, and enforcing a tighter preemption policy each raise one goodput factor by shifting loss to another boundary. Evaluating any cross-boundary change on the single factor it was built to improve ignores the capacity lost at the boundary that paid for it, so the change must be judged by the goodput product it leaves behind. This trade-off is the fleet-level manifestation of the displacement of overhead (principle 14), in which an optimization relocates overhead among compute, communication, and coordination rather than eliminating it outright. Because every factor multiplies the same goodput product, and one point of any factor is worth the product of the other two (section 9.8.1), a loss displaced to one boundary can be weighed in a common currency against the gain it purchases at another.

↪ Lifecycle cost of ownership, chargeback, and utilization as the invariant beneath build-versus-buy are developed in section 12.9.

9.9.1 Judging a trade by the product

When a compiler team enables an autotuning pass that speeds up steady-state steps while adding minutes to startup, the compiler team sees only the faster step trace and the reliability team sees only the longer launch in the job timeline. Deciding whether that trade helped the fleet requires combining both teams' measurements without rebuilding the whole-fleet waterfall from scratch. A change that touches more than one boundary multiplies MPG (equation 9.1) by the product of each factor's before-and-after ratio, with every ratio measured on the same boundary and over the same time window. The change improves fleet efficiency only when that product of ratios exceeds one, regardless of which individual factor the lever was designed to raise.

Expressing the rule in ratios allows separate engineering teams to evaluate a trade from their own telemetry: the reliability team measures what a restart policy does to runtime goodput, the compiler team measures what an autotuned build does to program goodput, and the two relative ratios multiply directly. Because a 1 percent relative gain in any factor produces a 1 percent gain in the overall product, small changes reduce to adding percentage gains and subtracting percentage losses. The marginal value of one absolute percentage point derived in section 9.8.1, the product of the other two factors, is the exact same rule expressed in percentage points. A point added to the lowest factor therefore represents the largest relative change and carries the highest marginal value.

Even so, every factor's ratio must be re-measured after the same change, because a lever aimed at one factor mechanically shifts the others across their shared boundaries. Productive step time is both the denominator of program goodput and the numerator of runtime goodput, while held allocation time is both the denominator of runtime goodput and the numerator of scheduling goodput. Shortening the time spent on one side of a boundary therefore alters the ratio on the other side. The fleet-goodput study measured this coupling on production clusters, and table 9.9 summarizes the directional effects (Wongpanich et al. 2025).

A compiler optimization that shortens on-device execution per step raises program goodput, yet it simultaneously lowers runtime goodput because the job now spends fewer productive device hours relative to its fixed startup, checkpoint, and recovery durations. When the job is device-bound, total wall-clock step time falls and the run finishes earlier. Releasing the allocation sooner shortens the job's held time, but the gap between one job's release and the next job's complete gang does not shrink with it. When each job holds its devices for less time, those fixed gaps take a slightly larger share of the pool's calendar window, so scheduling goodput dips slightly even as total MPG rises. Eliminating runtime waste likewise finishes the job earlier, raising runtime goodput while slightly lowering scheduling goodput, and only shortening partial-gang holding raises scheduling goodput in isolation.

When the job is host-bound and paced by its input pipeline, faster on-device execution does not shorten wall-clock step time at all. Every millisecond saved on the device becomes an extra millisecond spent waiting for the next input batch, so the rise in program goodput and the drop in

runtime goodput cancel out and leave MPG unchanged. Consequently, a monitoring dashboard that shows one factor dipping after an optimization does not indicate a regression until the full product has been evaluated.

Table 9.9: How a Lever Moves All Three Factors: Each row is a change made at one layer of the fleet stack and each column a factor's direction afterward. Program and runtime goodput share the productive-time boundary, so a faster step that raises one lowers the other, and only the product says whether the fleet gained. Adapted from Wongpanich et al. (2025).

Lever	Program goodput	Runtime goodput	Scheduling goodput	MPG
Compiler shortens on-device step	Rises	Falls	Falls if device-bound; else no change	Rises if device-bound; else no change
Runtime removes idle or lost time	No change	Rises	Falls	Rises
Scheduler shortens partial holding	No change	No change	Rises	Rises

This ratio rule is strongly asymmetric when the factors start at different levels. A program optimization that lifted the running example's program goodput from 0.48 to its single-node ceiling of 0.65 would increase the product by 35 percent. It would therefore still yield a net gain even if it sacrificed up to 26 percent of the job's runtime goodput (figure 9.13). A runtime lever has no comparable margin to spend: even if every runtime loss were eliminated, a runtime goodput of 0.90 could rise by at most 11 percent. Because a factor already near one has little relative headroom left to offer, a trade that spends runtime goodput to purchase program goodput has a wide margin before it lowers the product, although it still needs measuring, whereas a trade in the opposite direction has little margin and warrants immediate skepticism.

Two conditions determine whose product and which measurement window the product rule requires. First, a lever confined to a single job's allocation—such as a compiler flag or a local checkpoint strategy—is measured on that job's own product ($\eta_{\text{runtime}} \times \eta_{\text{program}}$). By contrast, a lever that moves capacity between jobs—such as priority preemption, a placement wait, or an elastic resize that returns devices to the pool—also alters other jobs' factors. It must therefore be judged on the pool's MPG, the only quantity that captures both the urgent job's scheduling gain and the preempted victim's runtime loss. Second, one-time costs such as a compilation build, a mesh rescale, or a restart purchase gains across every subsequent step, so a ratio measured over one hour and the same ratio measured over one week can fall on opposite sides of unity. For that reason, the measurement contract of section 9.3.1 requires the observation window to be reported alongside the result.

The goodput product itself does not measure wall-clock completion time: two configurations with the exact same product can finish a job days apart. The product rule evaluates capacity efficiency for the fleet, whereas section 9.9.4 introduces the additional dimensions on which a job's owner chooses an operating point.

The elastic scaling decision analyzed in section 8.7.3 illustrates this trade-off across boundaries. Consider an elastic job, separate from the running example, whose target batch is designed for 512 GPUs under three strategies. The first starts immediately on 128 GPUs and remains at that size. The second waits 4 hours in the queue for the full 512 GPUs allocation. The last starts immediately on 128 GPUs and steps up to 256 GPUs when capacity frees, leaving the remaining devices to queued jobs. Because the global batch size is held fixed across all three strategies, each epoch delivers identical statistical progress. Program goodput separates the three options most strongly: the scaling efficiencies—90 percent at 256 GPUs and 80 percent at 512 GPUs—are the exact ratios of each larger allocation's program goodput to that of the 128 GPUs baseline. Runtime goodput separates them by a single non-productive interval: the 10-minute pause required to checkpoint, reshard, and restart during the step-up transition (section 9.5.1). The step-up strategy retains its throughput advantage over the immediate small run unless that rescale pause exceeds roughly 9 hours.

Scheduling goodput—which belongs to the shared pool—determines which strategy serves the cluster best. The 256 GPUs that the step-up strategy leaves in the pool are worth whatever the queue can do with them. If waiting jobs form complete allocations on those devices immediately, the pool's

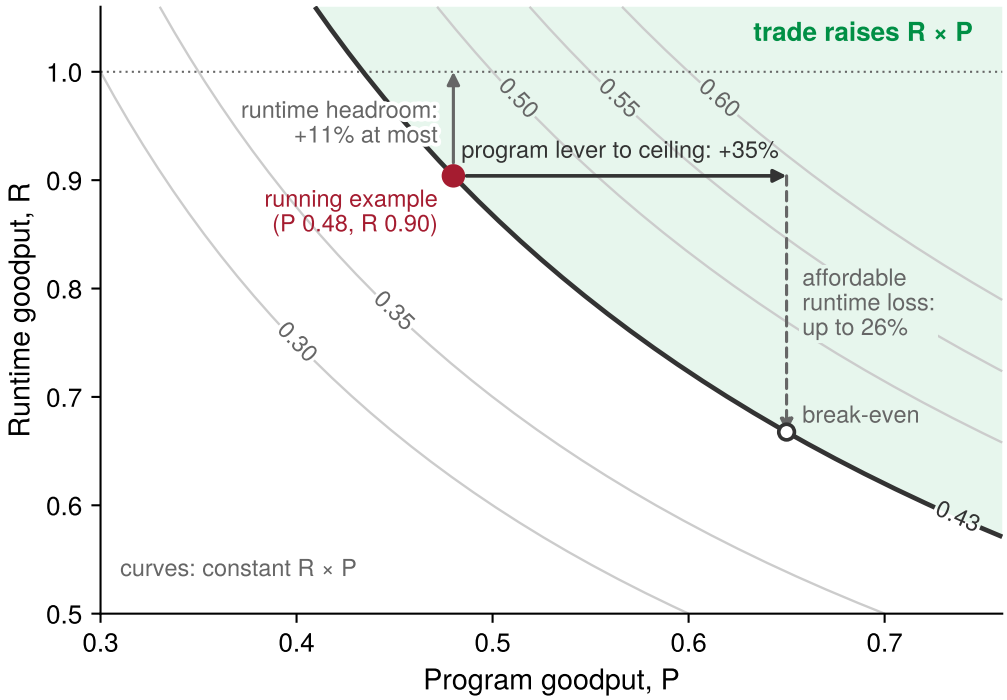

Figure 9.13: Iso-goodput Curves of the Product Rule: Program goodput against runtime goodput for the running example, with scheduling goodput held at 0.65, so each curve of constant $\eta_{\text{runtime}} \times \eta_{\text{program}}$ is also a curve of constant MPG. The example's P comes from constructed microbatch-scale step times, is illustrative, and is the per-microbatch value; amortizing the data-axis gradient reduction over the accumulated microbatches lowers it to about 0.456 over a full optimizer step and MPG from about 0.282 to about 0.268. A program lever that lifts P from 0.48 to its ceiling of 0.65 can give up as much as 26 percent of runtime goodput before it falls back to the bold curve through the running example, while removing every runtime loss raises R by only 11 percent.

MPG rises. If those devices sit in unallocated fragments that no queued gang can use, they become a scheduling loss, and giving the full allocation to the first job would have been more productive. Wall-clock completion time pulls in the opposite direction: the step-up strategy advances at nearly 1.7× the rate of the immediate small run yet still trails the full allocation by roughly 38 percent over the window. Consequently, a team facing a hard deadline prefers waiting for the full allocation, whereas a fleet operator managing a deep queue prefers the elastic step-up.

9.9.2 Cross-boundary trades across the fleet stack

Elastic resizing is one instance of a recurring pattern across the fleet stack: replacing a straggling node, preempting a low-priority job, lending idle quota, or staging checkpoints asynchronously in background memory each buys a gain at one goodput boundary by paying a cost at another. Earlier chapters in this book priced these mechanisms in the local units of their own subsystems, such as restart minutes, cascade multipliers, or storage bandwidth. Mapping those results onto the three goodput boundaries identifies which factor each lever raises and which factor it lowers, allowing cross-layer trade-offs to be compared on a single scale (table 9.10).

Table 9.10: Trades between Goodput Factors Priced in This Book: Each row reads an owning chapter's result onto the accounting boundaries, naming the factor the lever raises and the factor it lowers. The owners' results are not re-derived; the product rule of section 9.9.1 decides each trade from the measured ratios. The last row trades within one factor and is listed because it couples to the preemption rows.

Trade	Raises	Lowers	Priced in	What decides it
Straggler replacement	Program goodput: steps return to the healthy pace	Runtime goodput: restart, state reload, and replay of lost steps; any wait for spare capacity	section 7.9.5	How long the slowdown persists, against the restart and replay cost
Priority preemption	Scheduling goodput: the urgent gang forms without draining capacity for it	Runtime goodput of every victim, multiplied by the cascade	section 8.8.1	Preemption budget and requeue rule, against admission delay
Quota lending	Scheduling goodput: unused guaranteed capacity becomes a borrower's complete allocation	Runtime goodput of the borrower when the owner reclaims	section 8.11.2	How quickly a borrower can checkpoint or shrink
Placement wait	Program goodput: collectives stay on the fastest links	Scheduling goodput while capacity is held free to assemble the placement; otherwise only time to result	section 8.6.2	End-to-end completion time, including the wait
Elastic resizing	Program goodput: a smaller allocation runs at higher scaling efficiency; scheduling goodput when freed devices form other jobs' allocations	Runtime goodput: each rescale pause; the job's own time to result	section 8.7.3; section 9.9.1	The pool's queue and fragmentation at the resize
Asynchronous checkpoint staging	Runtime goodput: no synchronous pause per save	Program goodput: background writes contend with steps; runtime goodput again if a failure lands before the new generation is durable	section 4.6.1; section 7.8.3.1	Not priced in this book; measured step time with writes in flight
Compilation build	Program goodput: faster compiled steps	Runtime goodput: build time in every startup that compiles	section 9.9.3	Number of executions that reuse the artifact
Checkpoint cadence (within one factor)	Runtime goodput: less rework per interruption	Runtime goodput: more frequent pauses	equation 7.8; section 7.8.1; section 9.5.2	Interruption rate, including preemptions

Three structural patterns recur across table 9.10. First, any lever that transfers capacity between jobs—such as priority preemption or quota lending—must be evaluated at the pool level because its gain lands on one job while its cost lands on another. The orchestrator's preemption budget sets the rate at which one is exchanged for the other. Second, waiting in the queue costs scheduling goodput only when capacity is held idle to assemble the target allocation. If other jobs run on those devices in the meantime, the wait costs only wall-clock time to result, which is why placement waits and elastic resizing depend on the state of the pool's queue rather than on the job alone. Third, an elastic resize changes the device mesh, forcing a compiler-partitioned program to recompile for the new topology (section 9.7.1) and adding that build time to the rescale pause.

Asynchronous checkpoint staging is the one trade in the table whose program-goodput cost must be measured empirically on the target system. Staging copies a state snapshot to host memory or local flash and flushes it to durable storage in the background while training continues (section 7.8.3.1), removing most of the synchronous pause counted by the checkpoint term of equation 9.4. However, the background flush still contends with active training steps for pinned host memory, PCIe transfers, network links, and storage bandwidth, so its overhead surfaces as slightly longer retained steps—a program goodput loss. Staging also delays the durability point: until every shard reaches persistent storage, a failure rolls the job back to the previous checkpoint generation, increasing runtime rework if a fault occurs during the flush. Whether asynchronous staging raises the overall product therefore

depends on two measurements under the contract: the retained-step duration with and without a background flush in flight, and the time required to commit a checkpoint generation.

By contrast, checkpoint cadence trades two losses within a single factor—the time spent writing checkpoints versus the steps lost when an interruption occurs—so the Young-Daly interval (equation 7.7) minimizes total runtime loss rather than balancing two separate goodput factors. It couples to the orchestration rows through the total interruption rate, which is why including budgeted preemptions shortened the running example's optimal checkpoint interval in section 9.5.2. Graceful preemption, which saves state upon preemption notice rather than discarding the steps since the last periodic checkpoint, similarly trades rework against a synchronous save within runtime goodput.

9.9.3 Compilation builds

An autotuning compiler that searches for many minutes before the first step shortens every subsequent step by fusing operations, planning memory layouts, and selecting faster kernels, yet the accelerators sit allocated without retaining a single step while that search runs. Compilation is the cross-factor trade owned directly by performance engineering because both of its effects land inside the job's own allocation. The faster compiled steps raise program goodput (section 9.7.3), whereas the build and warmup that precede them are held time without retained progress, which counts as startup and lowers runtime goodput (section 9.5.1). A compilation build pays off when the cumulative time it saves across all executions that reuse the compiled artifact exceeds the build time itself; the break-even point therefore depends on the number of executions that share the artifact rather than on wall-clock deployment duration alone.

The boundary rule determines where a recompilation is charged. A build that runs at job launch, after a failure restart, or after an elastic resize to a new mesh occurs once per launch and is charged to runtime goodput along with the rest of startup. Conversely, a recompilation triggered repeatedly during steady-state execution—because input shapes vary or the captured graph keeps breaking—recurs inside the job's fault-free step loop and lowers program goodput instead. A dynamic shape distribution that forces a recompile every few hundred steps therefore manifests as slow retained steps rather than as startup overhead, and its remedy is a shape-polymorphic compilation strategy rather than a faster compiler build.

Napkin Math 9.6: When a compilation build pays for itself

Problem: A 13B-parameter model is small enough that one H100 holds it, so one accelerator forms a complete replica, here serving INT4 weights. That replica decodes at 120 tokens/s without compilation, and an autotuning compilation of its graph takes about 15 min to build. How many generated tokens repay that build on one replica, and how does the answer change when 128 replicas share one compiled artifact versus when each replica builds its own?

Variables:

- The uncompiled step takes 8.33 ms per token, of which matrix multiplications account for 0.50, element-wise kernels account for 0.35, and exposed launch overhead accounts for 0.15.
- Compilation leaves the matrix multiplications unchanged, removes 70 percent of the element-wise time through fusion, and removes 80 percent of the launch overhead.
- A default-mode build takes about 1 min; the autotuning search takes about 15 min.

Math:

- Compiled step: $0.50 + 0.35 \times 0.30 + 0.15 \times 0.20 = 0.635$ of the uncompiled step, about 5.3 ms per token, a 1.58× speedup that saves about 3.0 ms per token.
- One replica: 15 min $\div$ 3.0 ms $\approx$ 300,000 tokens, about 26 min of single-stream generation. The default build would need only about 20,000 tokens if it captured the same saving.
- Shared artifact: 26 min $\div$ 128 replicas $\approx$ 12 seconds of the pool's generation.
- Per-replica builds: 128 replicas $\times$ 15 min $\approx$ 32 GPU-hours per rollout, about 25 percent of the held time of a replica added for one burst hour.

Systems insight: A compilation build is amortized over executions rather than over elapsed time, so the same autotuning search that the pool repays in seconds when its artifact is shared consumes a quarter of a short-lived replica's held time when every replica repeats it independently.

A stored compilation artifact—or the autotuning cache that produced it—is built once for each combination of model version, shape range, and accelerator type, and its step-time savings accrue across every replica that loads it. By contrast, a per-replica build imposes a runtime loss on every replica, sized against the replica's lifetime in the same manner as the weight load and warmup analyzed in section 9.5.3. Because replicas added for a traffic burst have the shortest lifetimes, per-replica compilation makes the exact capacity a fleet adds under urgency both slower to come online and less productive while it runs.

Training jobs face the same trade-off across relaunches rather than across serving replicas. Unless compiled artifacts are cached persistently, a training job recompiles after every failure restart, every preemption, and every elastic resize to a new mesh, embedding compilation overhead into three terms of equation 9.4 at once: initial startup, failure recovery warmup, and preemption warmup. Whereas startup is negligible for a week-long pretraining run, it represents a substantial fraction of held time for hour-long jobs (section 9.5.2)—and production fleets run large numbers of short jobs, including evaluations, debugging runs, and hyperparameter sweeps. A persistent compilation cache[9] shared across launches turns a recurring per-launch runtime loss into a single build per program revision. Because the cache is keyed by the model revision, input shapes, device mesh, accelerator model, and compiler version, a compiler or hardware upgrade invalidates cached artifacts and pays every build anew, a cost caught by the startup regression checks of table 9.4.

Production compilation stacks differ in build cost, in the shape variations a compiled artifact tolerates, and in how broadly artifacts can be reused across launches (table 9.11). Deeper hardware specialization pays off only when the model revision and its shape distribution remain stable long enough for many executions to amortize a single build.

Table 9.11: Compilation Trade-Offs: Deeper specialization raises build cost and narrows the shapes a compiled artifact accepts, so the right path depends on how long a model version and its shape distribution stay fixed and on how many jobs and replicas reuse the artifact. Build costs are orders of magnitude for a large model, not measurements.

Compilation path	Build cost	Shape flexibility	Specialization depth	Reused across jobs by	Typical fit
`torch.compile` (default)	Seconds to minutes	High; graph breaks and recompiles on new shapes	Element-wise fusion, memory planning	A cache keyed by graph, shapes, and software version	Development and changing models
`torch.compile` (autotune)	Minutes to tens of minutes	Moderate; tuned per shape bucket	Autotuned tiles and launch configurations	Stored tuning results per shape bucket and accelerator	Stable shapes with long replay
XLA (JAX)	Minutes per shape signature	Low; new shapes or dtypes respecialize	Whole-region fusion, layout, and sharding	A cache keyed by program, shape signature, and device mesh	Static-graph training and TPU serving
TensorRT engine	Tens of minutes for a large model	Low; declared shape profiles only	Per-kernel tactic search, calibrated precision	A serialized engine per shape profile, accelerator model, and library version	Stable, high-volume inference

Graph breaks and shape specialization express the same trade-off at a finer granularity. A graph break—where dynamic control flow or unsupported host operations prevent the compiler from tracing a single graph—partitions the regions over which the compiler can fuse kernels and plan memory, and it can trigger recompilation when surrounding code paths change. Similarly, specializing to a narrow shape profile yields faster kernels but triggers a rebuild whenever a new shape arrives, whereas symbolic dimensions accept variable batch sizes and sequence lengths without rebuilding. Under the boundary rule, a poorly matched specialization strategy shows up either as

9 **Persistent compilation caches**: JAX writes compiled executables to a persistent cache directory and keys each entry by the unoptimized program, the jaxlib version, the relevant compiler flags, and the device count and topology. PyTorch's compiler keeps on-disk graph and kernel caches that later processes reuse, and it accepts an entry only for the same PyTorch and Triton versions and the same GPU. TensorRT serializes a built engine so that it can be reloaded, and by default the engine runs only on the GPU model and TensorRT version that built it. In all three frameworks the key is the tuple the text names, so an upgrade to the compiler or the accelerator turns every cached build into a new one.

slow retained steps (when a narrowly specialized artifact recompiles repeatedly during steady-state operation) or as inflated startup overhead (when an expensive autotuning search runs anew at every launch).

9.9.4 The efficiency frontier

For the elastic job of section 9.9.1, holding the job to 128 GPUs rather than 512 gives the higher per-device program goodput, whereas expanding it to 512 GPUs lowers program goodput yet finishes the run in a fraction of the wall-clock time. Because a workload's owner cares when the result finishes and what it costs as well as how efficiently the cluster runs, a cross-factor trade that raises the goodput product is still only part of an operational decision. The broader engineering objective is to choose a position on the **efficiency frontier**: the Pareto-optimal subset of feasible configurations for a stated workload and resource budget on which no measured dimension can improve without degrading another. As scaling laws drive rapid growth in the computation required for model quality improvements (Thompson et al. 2021), fleet efficiency determines how much of that purchased compute becomes useful progress, making a target quality reachable sooner or at lower cost.

↩ The power-law relation between compute and model quality that drives fleet growth is established in section 1.3.

Three measured dimensions span the efficiency frontier of a training fleet. First, goodput is the pool's MPG—or the job's product $\eta_{\text{runtime}} \times \eta_{\text{program}}$—measured under the contract. Second, cost per useful unit is the raw capacity price divided by that goodput (equation 9.6). Third, time to result is the wall-clock duration from job submission until the workload reaches its target quality. With the job's queue wait W, its allocated device count N, and the ideal accelerator-time T_{ideal} required to complete the work at the ideal rate, the time to result T_{result} follows directly from the job's goodput factors (equation 9.7):

$$T_{\text{result}} = W + \frac{T_{\text{ideal}}}{N \times \eta_{\text{runtime}} \times \eta_{\text{program}}} \tag{9.7}$$

Scheduling goodput does not appear explicitly in the denominator of equation 9.7 because unallocated capacity elsewhere in the pool does not slow a job that already holds its gang. The pool's contention enters the job's completion time solely through the initial queue wait W. The allocated device count N does appear in the denominator, but program goodput η_{program} declines as N grows according to the layout's scaling efficiency. Adding devices therefore shortens held execution time only while the product $N \times \eta_{\text{program}}$ continues to rise, and each additional device lowers per-device program goodput. Consequently, a deadline-driven job purchases shorter completion time by holding more devices at lower program goodput, whereas a heavily loaded fleet maximizes total useful output by holding that job to a smaller, more efficient allocation. The three elastic scaling strategies analyzed in section 9.9.1 occupy three distinct points along this trade-off curve, and none of them Pareto-dominates the others.

On a homogeneous fleet with a single capacity price, goodput and cost per useful unit are reciprocals of each other, reducing the frontier to a two-dimensional trade-off between goodput and time to result. The two dimensions separate whenever capacity tiers carry different prices. Preemptible spot capacity is discounted relative to reserved capacity but experiences more frequent interruptions, lowering runtime goodput alongside the hourly price; the spot-cost equation (equation 8.3) determines whether the price discount outweighs the runtime rework. Under the elastic scaling parameters of section 8.7.3 and the spot interruption profile of section 8.13, the price discount more than compensates for the interruptions, placing spot configurations on every point of the cost-time frontier except the single fastest completion time (figure 9.14). Choosing between reserved and on-demand instances, retaining older accelerator generations at a lower internal chargeback rate, or introducing a second accelerator family all pose the exact same trade-off, and each option must be evaluated by cost per useful unit rather than by raw price per accelerator-hour.

Model quality and device memory bound the set of feasible configurations rather than ranking the points within it (table 9.12). A configuration that runs faster by lowering numerical precision, enlarging the global batch size, or relaxing gradient synchronization is valid only if it still reaches the measurement contract's quality target, which is why time to result is always measured to a fixed quality criterion (section 9.10). Similarly, a parallelism layout, microbatch size, or sequence length that exceeds device memory does not lie on the feasible frontier at all. The mechanisms that make it

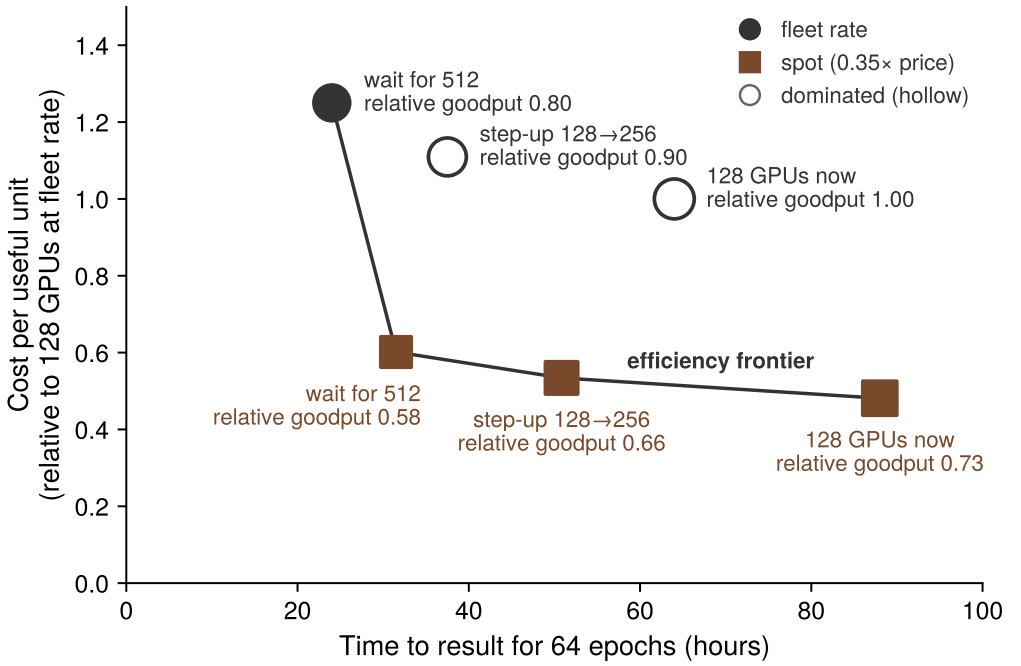

Figure 9.14: Efficiency Frontier of the Elastic Strategies at Two Prices: Time to result and cost per useful unit for one fixed target, the 64 epochs the full allocation completes in the 24-hour window of the elastic scaling decision in section 8.7.3, under its three strategies on the fleet rate and on spot capacity at 35 percent of that rate; marker area and the labels' relative goodput give each strategy's useful work per held GPU-hour as a ratio to the 128-GPU run on the fleet rate (1.00), not the running example's absolute goodput. At one price, cost per useful unit is the price divided by goodput, so cost and goodput move together and the three fleet-rate points trade time against cost with none dominating another. Spot's discount outlasts its interruption losses here, so the frontier runs through every spot point and only the fastest fleet-rate point; the other two fall behind it (hollow). Values are illustrative, derived from that decision's inputs and the spot interruption scenario of section 8.13; spot points keep the fleet-rate queue wait and count each interruption as a graceful save.

fit, such as recomputing activations or sharding model state more finely, purchase feasibility at the expense of program goodput.

Table 9.12: Efficiency Frontier Dimensions: Fleet efficiency places a configuration by its goodput, its cost per useful unit, and its time to result, within the set that quality and memory make feasible. On one fleet at one price, goodput and cost per useful unit move together; they separate when options differ in price. Latency under offered load is a serving dimension and is developed with the serving system.

Dimension	How it is measured	What it constrains	Typical tension
Goodput	MPG of the pool, or runtime times program goodput of a job, under the measurement contract	How much of the capacity paid for becomes progress at the ideal rate	A lever that raises one factor often lowers another (table 9.10)
Cost per useful unit	Capacity price divided by goodput (equation 9.6)	What each unit of progress costs	Cheaper capacity, such as preemptible or older accelerators, often carries lower runtime or program goodput
Time to result	Queue wait plus held time to the quality target (equation 9.7)	When a job's result is available	More devices shorten held time but lower program goodput and can lengthen the queue wait
Quality	Accuracy or loss at the target, under the contract's quality criterion	Which output counts as useful work	Lower precision, larger batches, and relaxed synchronization need a quality check before their speed counts

Dimension	How it is measured	What it constrains	Typical tension
Memory	Peak device memory per rank	Which layouts, batch sizes, and sequence lengths are feasible	Recomputation and finer sharding make a configuration fit at the cost of program goodput
Latency under load (serving)	Request latency at a stated offered load	Whether a replica meets its service-level objective	Developed in chapter 10; enters this accounting as headroom charged to program goodput

Because different stakeholders value different points along the same frontier, an engineering proposal should state which frontier dimension it targets before being evaluated. A flagship training run facing a hard release deadline accepts a higher cost per useful unit to shorten its time to result, holding a larger allocation and waiting for a compact placement. A research sweep comprising hundreds of short jobs optimizes for cost per useful unit, favoring smaller allocations, preemptible spot capacity, and a shared compilation cache. The cluster infrastructure team optimizes for the pool's overall MPG and favors policies that raise the fleet-wide product even when they lengthen individual jobs. The product rule determines whether a change moves the system toward the efficiency frontier, while the stakeholder's operational priority determines where along that frontier to operate.

Any move along the frontier is verified using the same discipline as a single-step optimization. Engineers measure every affected goodput factor before and after the move under the measurement contract, over an observation window that captures both one-time setup costs and recurring per-step gains. They report the change in time to result alongside the ratio of the goodput product.

Checkpoint 9.4: Trading between factors

A lever that raises one goodput factor and lowers another is judged by the product of the factors' ratios, and the efficiency frontier places the result against cost and time to result.

Judging the trade

- ☐ **Ratios, not points**: A change raises program goodput by a tenth of its value and lowers runtime goodput by a twelfth of its value. Can you decide whether the product rose without knowing scheduling goodput, and explain why?
- ☐ **Whose product**: Why must a preemption be judged on the pool's product rather than on the product of the job it admits?
- ☐ **Within one factor**: Why is the checkpoint interval a trade inside runtime goodput, and how does a preemption budget change it?

Compilation and the frontier

- ☐ **Compile amortization**: When does a shared compiled artifact repay its build, and when does per-replica or per-launch compilation lower runtime goodput more than the compiled speed raises program goodput?
- ☐ **Time and goodput**: Using the time-to-result equation, explain why holding a job to fewer devices can raise the fleet's goodput while lengthening that job's time to result.

9.10 Fallacies and Pitfalls

An organization can report a fleet MFU above the best published flagship runs while its shared cluster turns little more than a quarter of its offered capacity into training progress. Each misconception below treats a single factor, a single observation window, or a single kernel as the entire efficiency story; the corresponding goodput boundary exposes what that narrow view leaves out.

Fallacy: *A job with high MFU is using the fleet well.*

MFU measures program goodput and nothing else. Because it divides only by the time spent inside retained steps, it is blind to capacity that the job never held and to held time that produced no retained progress. In our running example, retained steps execute at a program goodput of 0.48—above the 0.41 reported for the Llama 3 405B pretraining run—yet the pool's scheduling goodput of 0.65 and the job's runtime goodput of 0.90 leave a pool-level MPG of only 0.28. Only the full product of equation 9.1, composed in section 9.8.1, states how much of the fleet's capacity becomes useful work.

Fallacy: *A fleet with high capacity, occupancy, and duty cycle is productive.*

The three traditional cluster metrics count whether accelerators are available, allocated, or electrically active, but none of them counts useful model progress. Availability counts free accelerators without regard to network topology, so a distributed job requiring a contiguous mesh can sit blocked in the queue beside a large but fragmented pool. Occupancy counts accelerators assigned to jobs, so a device stalled on input data, executing unoptimized kernels, or repeating steps lost to preemption still registers as fully occupied. Duty cycle counts whether a device's execution engines are active, not what fraction of peak throughput they achieve or whether their operations advance the model. The fleet-goodput framework replaces all three metrics with the goodput factors precisely because each factor tracks capacity-time across the next boundary of useful progress rather than stopping as soon as the hardware appears busy (Wongpanich et al. 2025).

Pitfall: *Charging each loss to the team that will fix it.*

Assigning losses by engineering team routes each problem to someone who can act on it, but doing so inside the goodput factors breaks the identity that allows those factors to multiply. Because a zombie job's accelerators are already counted as held in the scheduling numerator, charging zombie time to scheduling goodput—on the grounds that the orchestrator's lease mechanism will clean it up—subtracts the same lost hours twice. Similarly, charging a slow straggler node to reliability makes the slowdown vanish from program goodput, even though every synchronous step runs at the pace of that slowest rank. To preserve multiplicative accounting, the boundary where time is lost determines which factor records the loss, while the root cause determines which engineering team owns the fix (section 9.2.3); a complete attribution table therefore carries both columns side by side (table 9.5).

Fallacy: *A single profiling run is sufficient to characterize a job's goodput.*

A single short run can characterize at most one factor, because the three goodput factors lose time on timescales that differ by orders of magnitude. Program goodput loses time inside every step, so a window of several hundred warmed steady-state steps estimates it reliably. Runtime goodput, by contrast, loses time in rare, expensive interruptions: on 128 GPUs, the running example encounters a hardware fault roughly once every 13 days, meaning a single-day profile typically captures no failure and no recovery. Scheduling goodput varies with cluster-wide demand, which fills queues during peak hours and drains overnight. For this reason, the measurement contract (section 9.3.1) specifies both the factor and its required observation window, and fleet sampling (section 9.3.3) spreads short program traces across that longer window.

Pitfall: *Subtracting one overhead budget that mixes factors.*

A spreadsheet planning budget that reserves a few percent each for pipeline bubbles, checkpoints, failure recovery, scheduling delay, and preemption—and then subtracts their sum from one operational ratio—conflates all three goodput factors. Pipeline bubbles occur inside retained steps and belong to program goodput; scheduling delay precedes gang allocation and belongs to scheduling goodput; only checkpoints, failure recovery, and preemption belong to runtime goodput. Multiplying the remainder of such a mixed budget by a measured fleet MFU double-counts pipeline bubbles and folds scheduling losses into a factor that has no scheduling term. A sound capacity budget assigns each loss to its proper boundary first and sums only the terms that share the same factor (section 9.5.2).

Fallacy: *More FLOP/s means a proportionally faster step.*

Raising peak arithmetic throughput shortens only the compute-bound portion of the iron law, and only where that computation is exposed on the critical path. The serving replica's batch-one decode step uses roughly 0.3 percent of its FP16 arithmetic peak; doubling peak FLOP/s at the same HBM bandwidth leaves its 5.2 ms latency floor untouched and merely cuts its numerical MFU in half. Likewise, the running training example spends 65 ms of each 245 ms illustrative fleet step above its single-node baseline, in what the example assumes is cross-node synchronization and collective waits that faster arithmetic does not accelerate, and adopting a lower-precision number format encounters the exact same Amdahl limit (section 9.7.3). Evaluating a step against the roofline model and the Fleet Law terms (section 9.6.3) predicts how much speedup a faster accelerator will actually deliver.

Pitfall: *Judging a job by its mean step time while ignoring the step-time distribution.*

A healthy mean step time can mask severe tail episodes. Because a synchronous distributed step completes only when its slowest rank finishes, a single device that experiences intermittent slowdowns—from thermal throttling, link bit-error retries, or host OS pauses—delays its entire collective group on every episode. Averaging across thousands of steps smears those stalls into what looks like a minor, uniform slowdown. Per-rank step-time percentiles expose those episodes as tail outliers and trigger targeted traces comparing the slow rank against a healthy peer (section 9.3.3), while the gray-failure regression check of table 9.4 monitors the same tail ratio across the fleet. Every such tail stall lowers program goodput regardless of which subsystem team eventually replaces the degraded node.

Fallacy: *The largest kernel is where program goodput is lost.*

The largest kernels in a transformer step are typically its dense matrix multiplications, and in a well-tuned implementation they already execute near the accelerator's arithmetic ceiling. Normalization, activation, softmax, and optimizer kernels sit on the memory-bandwidth slope and perform a tiny fraction of the arithmetic, yet together they can consume a substantial share of on-device step time. At fleet scale, exposed inter-node collectives that do not appear in a single-GPU kernel ranking can outweigh them all. Indeed, the constructed drill-down for our running example (section 9.6.3) places its largest program loss in exposed cross-node synchronization and collective waits. Partitioning step time into matrix multiplications versus everything else, and matching the remainder against table 9.7, identifies which Fleet Law term holds the recoverable loss before any individual kernel is tuned.

Pitfall: *Treating a successful compilation as proof that the program is well optimized.*

A compiler applies only the transformations its passes implement under its internal cost model, and a completed build does not guarantee that those passes succeeded. Automatic sharding propagation inserts a resharding collective wherever two inferred tensor layouts conflict, allowing an unintended AllGather to appear inside every transformer layer (section 9.7.1). Likewise, a dynamic shape distribution that triggers recompilation every few hundred steps surfaces as slow retained steps—lowering program goodput—rather than as a one-time startup cost (section 9.9.3). A compiled step should therefore be treated as a testable hypothesis, verified by checking the profiler trace's collective count against the intended layout and by measuring program goodput before and after under the measurement contract.

Fallacy: *Fixing the largest loss first recovers the most useful work.*

The loss shares of equation 9.2 describe where fleet capacity went in the past, not how much useful work an engineering lever will recover in the future. In the running example, scheduling loses 35 percent of fleet capacity and program execution loses 31 percent, yet a hypothetical perfect scheduler would convert only about 15 percent of fleet capacity into additional useful progress. Capacity recovered at the scheduling boundary must still pass through the runtime and program boundaries at their current rates. At the subsystem ceilings, the scheduling lever raises MPG from 0.28 to 0.35, whereas the program lever raises it to 0.38. Engineering priority is governed by the marginal value of one point of each factor—the product of the other two factors—and by each lever's reachable ceiling (section 9.8.1), not by the raw size of the historical loss slice.

Pitfall: *Counting retained steps as useful work without checking quality.*
MPG counts retained steps at the ideal rate under the assumption that each step advances model quality as much as it did before an optimization. A change that accelerates steps by lowering numerical precision, inflating the global batch size, or relaxing synchronization can raise program goodput while requiring more steps to reach the target loss, causing the fleet to report higher goodput while delivering the trained model later. Both the quality gate of the measurement contract (section 9.3.1) and the efficiency frontier's time-to-quality-target metric (section 9.9.4) catch this degradation, and statistical goodput (section 9.2.4) quantifies it directly. A throughput speedup counts as useful work only after the quality gate passes.

Each of these pitfalls measures a single boundary in isolation and mistakes it for the whole system. Returning every diagnosis and every trade-off to the full multiplicative accounting prevents local optimizations from masking fleet-level losses.

9.11 Summary

A machine learning fleet can be fully allocated and electrically busy while most of its capacity never becomes useful model progress. Capacity leaks away at three distinct boundaries across a workload's lifecycle: before a job holds a complete gang allocation, while the job holds accelerators without retaining progress, and inside retained steps that execute below the ideal rate. Scheduling goodput, runtime goodput, and program goodput measure what survives each boundary, and their product defines ML productivity goodput (equation 9.1). This multiplicative identity holds without double-counting only when every lost interval is charged once, at the first boundary it fails to cross.

Because the three factors lose time on different timescales—program losses inside every step, runtime losses during intermittent interruptions, and scheduling losses across diurnal demand cycles—a reproducible measurement contract must fix the observation window alongside the workload, the system state, and the timing boundary.

When composed from the owners' modeled terms, the running example's runtime loss came primarily from input stalls and preemptions rather than from hardware faults, demonstrating that identical headline factor values can point to entirely different subsystem owners. Across the shared pool, the three factors gave our running example an MPG of 0.28 and a cost per useful accelerator-hour equal to 3.5× the raw capacity price. Although scheduling lost the largest share of the pool's total capacity and program execution lost the largest share of the job's held time, the program boundary offered the highest-leverage engineering return. Its single largest loss, the exposed cross-node communication of each step, lay beyond the reach of any single-GPU kernel optimization.

Key Takeaways: Fleet efficiency is a product of owned factors

- **Useful work has three factors, each counted at one boundary**: Scheduling, runtime, and program goodput multiply into the share of fleet capacity that becomes useful work. Record each loss in the single factor whose boundary it fails to cross, and record the subsystem whose lever removes it in a separate attribution column, because boundary and owner rarely coincide.
- **MFU measures only program goodput**: High MFU proves that a job executes its retained steps efficiently, not that the fleet is used well. The running example's retained steps run at 0.48 of the ideal rate while only 28 percent of the pool's capacity becomes useful work, and bandwidth-bound decode workloads must be judged by MBU against a bandwidth floor instead.
- **Localize the factor before the mechanism**: Inspect scheduler logs first, job timelines second, and step traces third. Profiling a single step before localizing the lost factor can pronounce a job healthy while its capacity drains away in queue fragmentation or restart cascades that no step trace ever captures.
- **Levers act only through the term they shorten**: Kernel optimizations shorten local execution, sharding layouts set the communication volume, and overlap hides part of

that transfer time; consequently, any single-node kernel speedup is diluted at fleet scale as exposed communication and synchronization claim a larger share of the step.

- **Speedups do not add, so re-measure**: Shortening one exposed interval frequently uncovers the next bottleneck, and accelerating the backward pass can expose collective transfers that computation previously hid. Re-measure the distributed step under the same contract after every change.
- **A loss share is not a recoverable gain**: Capacity recovered at an early boundary must still pass through later boundaries, so one percentage point of any factor is worth the product of the other two. Rank engineering levers by their marginal product value and their reachable subsystem ceilings rather than by raw loss shares.
- **Judge trades by the product**: A lever that raises one factor while lowering another improves fleet efficiency only when the product of the factors' before-and-after ratios exceeds unity. Because adjacent factors share boundaries, a faster step raises program goodput while mechanically lowering runtime goodput, meaning a dip in one factor is not a regression if the product rises. A compilation build pays off only across the executions that reuse its cached artifact, and wall-clock time to result remains a distinct operational dimension on the efficiency frontier.

Ultimately, the cross-layer trade-offs in this chapter express the displacement of overhead (principle 14) in a single quantitative currency. Inside a training step, a faster compute kernel shifts exposure from local execution onto inter-node communication. Across a job's lifecycle, waiting for a compact placement shifts loss from program execution into the scheduler queue, an autotuning compilation shifts loss from step execution into startup, and maintaining a warm pool of serving replicas shifts cold-start latency into idle headroom. Overhead relocates across boundaries rather than vanishing, and the goodput product of the ML goodput decomposition (principle 16) determines, for every relocation, whether the fleet came out ahead.

What's Next: From useful work to served requests

Everything measured in this chapter averaged capacity over windows of hours to weeks and assumed that batches were waiting whenever an accelerator became ready. Serving at scale removes that assumption: user requests arrive stochastically and each must complete within a strict latency SLO, so queueing dynamics, continuous batching, and KV-cache admission determine how much of a replica's measured capacity becomes compliant responses. The serving replica analyzed here, characterized by its HBM bandwidth floor and its cold-start readiness time, provides the foundation for that analysis, and its governing metric, SLO goodput, filters capacity after MPG by counting only the requests that meet their latency targets.

↪ Request admission, batching, cache management, and the capacity a latency target requires are developed in chapter 10.

10

Serving at Scale

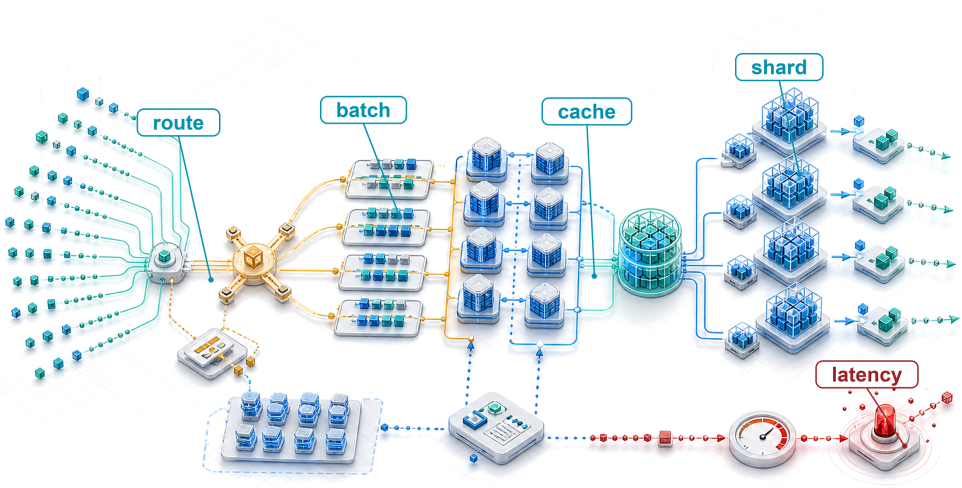

10.1 Inference Cost and Serving Architecture
10.2 Limits of Single-Machine Serving
10.3 SLO Goodput and the Serving Tax
10.4 Serving Cost Dominance in Production
10.5 Serving Workloads and the Serving Hierarchy
10.6 Batching Strategies at Scale
10.7 Memory and Decode-Time Management
10.8 Prefill and Decode
10.9 Weight Quantization for Serving
10.10 Model Sharding for Inference
10.11 Load Balancing and Request Routing
10.12 Multi-Tenancy and Isolation
10.13 Autoscaling and Capacity Control
10.14 Global Inference Infrastructure
10.15 Case Studies
10.16 Fallacies and Pitfalls
10.17 Summary

Purpose

Why can lifetime inference cost dwarf training cost?

Training is paid for each model version; serving continues throughout that version's operational lifetime. A model trained for millions of dollars may serve billions of requests, each consuming compute, memory, network, and energy. At high traffic, recurring inference expense can exceed the original training cost by orders of magnitude. Milliseconds of latency and small differences in accelerator utilization then compound across the service's lifetime. A fast prediction measured on an otherwise idle machine does not reveal how the same model behaves when requests arrive together, vary in length, and compete for memory. Work can spend longer waiting than executing, while an apparently busy fleet produces responses that arrive too late to be useful. Keeping extra machines ready protects response time but incurs cost during quiet periods; admitting more work improves utilization only while the service can still meet its promises. Failures sharpen this tension because displaced requests consume the spare capacity intended to absorb demand changes. Continuing to accept work during overload can worsen the shortage when delayed requests trigger retries that compete for capacity. The engineering decision is therefore how much useful work the fleet can sustain at an acceptable cost while preserving response quality and timeliness. The answer depends on the request mix and service requirements, so a gain for one workload may become a loss for another. In C^3 terms, compute reuse must pay for communication and coordination without exhausting the user's latency budget.

Learning Objectives

- Estimate lifetime serving cost from traffic volume, deployment duration, and measured cost per request that meets its SLO
- Select batching, scheduling, and per-request compute-budget policies for vision, LLM, recommender, and streaming workloads
- Analyze attention-cache capacity, fragmentation, and prefill-decode interference under workload-dependent serving constraints
- Compare sharding, disaggregation, and quantized serving designs using communication, memory, and quality constraints
- Evaluate routing, overload-control, and isolation policies against tail latency, retry amplification, and noisy-neighbor risks
- Design autoscaling and multi-region routing policies that balance cold starts, cost, and SLO compliance
- Synthesize serving architectures that manage C^3 trade-offs across request, replica, service, and platform layers

10.1 Inference Cost and Serving Architecture

A language model that costs millions of dollars to train over two months can, once deployed to a global user base at thousands of queries per second, consume a comparable sum in inference compute within weeks. Inference is then no longer a temporary batch job. It is a continuous, latency-sensitive service.

Running that continuous service takes more than an efficient steady-state replica. Fleet efficiency characterizes a replica by three quantities, each averaged over windows of hours to weeks. The first is model bandwidth utilization (MBU), which measures how close the replica's decode steps come to the bandwidth floor, the time needed to stream each step's weight and KV-cache bytes at peak memory bandwidth. The second is how long the replica takes to become ready. The third is the fraction of the time the replica holds its accelerators that it spends doing useful work at the ideal rate the bandwidth floor sets, which combines the runtime and program factors of ML goodput. Serving policy determines how that capacity is shared when requests arrive unevenly and each carries a latency target. A faster decode step supplies a better service rate, but it does not determine how long requests wait to join a batch, which replica receives them, which requests receive scarce cache space, or when the fleet needs another replica. These decisions connect the local execution toolkit to user-visible performance under load.

↩ The serving replica's program and runtime goodput, its bandwidth floor, and the fusion, graph replay, and compilation levers that shape it are established in chapter 9.

The service contract makes that distinction measurable. Offered load is the work arriving at the endpoint, including requests the service may reject; admitted load is the work it accepts. A service-level objective (SLO) specifies the required outcome over a measurement window, such as a latency percentile or successful-request fraction. For language models, request rate alone is insufficient because prompt and output lengths determine different amounts of work. Capacity comparisons must therefore hold the request mix and SLO fixed while measuring how much offered work completes successfully within the targets.

An admitted request passes through routing, preprocessing, queueing, batch formation, model execution, and response delivery, crossing between the host CPU and the accelerator twice on the way (figure 10.1). The host manages admission, tokenization, and detokenization, while the GPU executes the prompt prefill pass and iterative KV-cache decode loop. The figure maps noncompute serving overhead against the time-to-first-token (TTFT) and time-per-output-token (TPOT, the average gap between streamed tokens) latency budgets. For an autoregressive model, prefill processes the prompt and builds cached attention keys and values; decode extends that KV cache as it generates subsequent tokens. Waiting before prefill delays the first token, while competition during decode delays the stream. The scheduler must budget both delays alongside execution time. Batching and cache management govern work within a replica, sharding changes the replica's capacity, and routing and autoscaling distribute demand across replicas.

↩ Prefill's arithmetic-bound regime and decode's bandwidth floor, measured as MBU, are established in section 9.2.2.

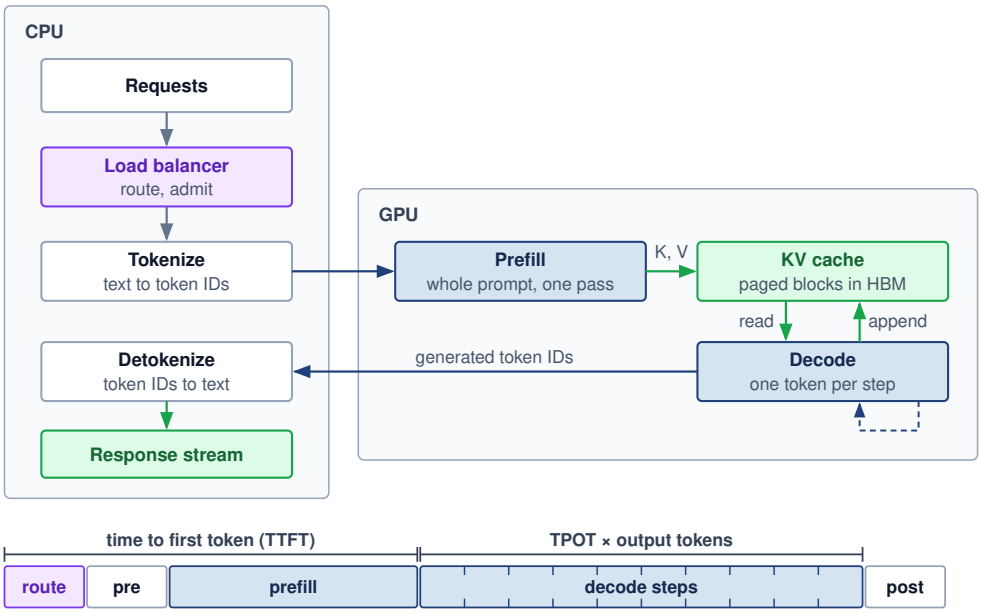

Figure 10.1: End-to-End Inference Pipeline: One language-model request crosses between CPU and GPU twice. The load balancer routes and admits it and the CPU tokenizes it; prefill then processes the whole prompt in one pass and writes its keys and values to the KV cache, and decode reads and appends that cache one token per step before the CPU detokenizes the stream. The strip orders one request's latency, with illustrative segment widths. Everything up to the end of prefill counts against time to first token, and the routing and CPU segments are serving-tax work that spends latency budget outside GPU compute.

Distribution adds placement and coordination decisions to the local serving problem. Load balancing[1] must distribute requests across replicas with different amounts of unfinished work and available memory. Request routing must also account for model-specific characteristics: recommendation systems with large sharded embedding tables require different placement strategies than language models that generate responses token by token. When demand changes faster than new replicas can become ready, routing alone cannot restore capacity. Autoscaling, spare capacity, and admission limits must work together to protect the service's latency targets.

Training and serving place costs on different timelines. A model version's training expense can be amortized over its deployment, while serving capacity must follow recurring demand. An e-commerce recommendation service, for example, must accommodate peak shopping traffic without paying for an unnecessarily large idle fleet during quiet periods. Both training efficiency and serving efficiency matter, but the serving decision must also account for when requests arrive and how long they can wait. Building such a service requires three linked decisions: when distribution becomes necessary, how the architecture meets latency targets under varying load, and how resource utilization stays high enough that serving cost does not erode the value the model creates. The worked scenarios that follow size these decisions using registry values for real hardware and models alongside labeled scenario assumptions for traffic, prices, and latency budgets, starting with the physical limits that force a service beyond a single machine.

[1] **Load balancing for inference**: Request service times can span milliseconds to tens of seconds as models, prompt lengths, and generated outputs vary. Queue-aware routing exposes that remaining work more directly than assignment counts alone.

10.2 Limits of Single-Machine Serving

A single accelerator fails a production endpoint in three ways: its memory fills before the model weights or caches fit, its throughput saturates below the arriving query rate, or its forward pass runs longer than a request's latency budget. Overcoming those limits requires distributing work across several devices within one machine or across machines over the network. Table 10.1 categorizes these three triggers and their matching distribution strategies, whose necessity depends on the permitted precision, placement, and service contract.

The first signal is memory exhaustion, which occurs when model parameters, key-value caches, or embedding tables exceed single-device capacity. A single NVIDIA H100 GPU provides 80 GB of HBM3[2] memory, whereas an illustrative mixture-of-experts model with ~1.8T parameters requires ~3.5 TB for FP16 weights alone. Keeping all those weights resident on accelerators requires multiple devices regardless of request volume; quantization and offloading change that placement constraint. Recommendation systems with trillion-parameter embedding tables face similar constraints. Meta's DLRM architecture[3] stores embedding tables that require multiple terabytes of memory (section G.2 records the provenance of the hardware and model constants used throughout this chapter).

Beyond memory constraints, throughput limitations emerge when request volume exceeds single-machine capacity even with optimal batching. Consider a recommendation system serving 100,000 queries per second with a 10 ms latency budget. If single-machine throughput peaks at 10,000 QPS, no amount of optimization on that machine can satisfy demand. Horizontal scaling across multiple replicas becomes mandatory.

Finally, latency requirements can motivate distribution even when memory is sufficient. The FP16 weights of a 70-billion parameter model, approximately 140 GB, exceed a single 80 GB H100-class GPU before KV cache is considered, but quantizing the weights to fewer bits or choosing larger-memory hardware can fit the model on one device. Even when the model fits, that device's execution time may still exceed the request budget. Sharding the model across multiple GPUs can divide computation and weight traffic, but it adds communication and synchronization. Because sharding shortens compute-bound prefill and bandwidth-bound decode by different amounts, the latency benefit must be measured for each phase against its own target.

[2] **HBM3 capacity and bandwidth**: The H100 combines 80 GB of 3D-stacked DRAM with 3.35 TB/s of bandwidth, compared with 2.04 TB/s for HBM2e on the A100. Capacity determines whether model and cache state fit; bandwidth bounds memory-bound decode throughput.

[3] **DLRM (Deep Learning Recommendation Model)**: Meta's 2019 reference architecture combines dense features processed by multilayer perceptrons (MLPs) with sparse features represented by embedding lookups (Naumov et al. 2019). This separation permits different placement strategies; it does not require embeddings on CPUs and MLPs on GPUs. Production recommendation characterizations show why this structure matters for serving: embedding-heavy recommendation models create memory-access and capacity constraints that differ from dense convolutional neural network (CNN), recurrent neural network (RNN), or large language model (LLM) inference (Gupta et al. 2020).

Table 10.1: Triggers for Distributed Inference: Each constraint type indicates different distribution strategies. Memory constraints require model sharding; throughput constraints require replication; latency constraints may require either depending on whether the bottleneck is compute or memory bandwidth.

Constraint	Single-machine limit	Example workload	Distribution strategy
Memory	80 GB (H100)	Illustrative MoE (~3.5 TB FP16)	Tensor/pipeline parallelism
Throughput	~10K QPS (vision)	100K QPS RecSys	Horizontal replication
Latency	Forward pass > SLO	500 ms LLM TTFT	Model sharding

Memory, throughput, and latency triggers all become user-visible once an interactive endpoint starts streaming. TTFT determines when the user sees a response begin, while TPOT determines whether the rest of the response arrives at a readable pace. TPOT is the average of the time between tokens (TBT), the gap between successive streamed tokens, which many serving systems report as inter-token latency (ITL). A stream can meet its TPOT target while a single long TBT gap still stalls the reader, so average cadence is stated here as TPOT and per-gap tail targets as TBT. Protecting both first-token delay and stream cadence is why inference reverses many of the optimization priorities from training.

Checkpoint 10.1: Distribution strategy selection

Verify your understanding of when to move from single-machine to distributed inference:

- ☐ **Shard vs. replicate**: for a 13B model (26 GB weights) that fits on a single A100, select model sharding or horizontal replication when throughput requirements double.
- ☐ **Latency floor**: for a real-time speech-to-text service where TBT is the primary bottleneck, determine whether data-parallel replication reduces this latency.
- ☐ **Mandatory sharding**: explain why a 175B model requires multiple 80 GB GPUs when all FP16 weights must remain accelerator-resident, regardless of request volume.
- ☐ **Latency trigger**: identify when a latency target forces distribution and when added network, serialization, coordination, or queuing overhead can erase the benefit.

10.3 SLO Goodput and the Serving Tax

A replica that runs near its hardware limit still wastes its capacity if requests wait so long in queues, network hops, and host processing that responses arrive after the user's deadline. Because a served request counts as useful work only when it meets its latency and quality targets, every stage it crosses spends part of that contract's budget.

10.3.1 The inversion from training to inference

A training iteration that takes 2 seconds instead of the usual 1 second slows a run slightly, and faster steps later can still keep the job on schedule. An interactive request that takes 2 seconds against a 100 ms target fails its contract immediately, even if every subsequent request finishes on time. This contrast between training and inference centers on the unit of useful progress. Training evaluates throughput together with convergence over a run; interactive serving evaluates completed requests against quality and latency targets. Offline inference can also prioritize throughput, and training still loses progress to stragglers and stalls, so the distinction is which delays violate the workload's contract, not whether one activity cares about time and the other does not. That difference reorders priorities across six system aspects, from the primary metric onward (table 10.2).

Table 10.2: Training vs. Inference System Requirements: The fundamental inversion from throughput to latency optimization ripples through system design, state management, batch formation, and failure recovery.

Aspect	Distributed training	Distributed inference
Primary metric	Throughput (samples/hour)	Latency (p99 ms)
Acceptable variance	Hours	Milliseconds
State management	Checkpoints (periodic)	Session state (continuous)
Batch formation	Large, controlled	Request-driven, variable
Failure tolerance	Restart from checkpoint	Redirect if capacity and state allow
Cost structure	Fixed duration, variable rate	Variable duration, fixed SLO

Failure handling follows those different contracts. Training restores a checkpoint and repeats lost work, charging recovery time to the run. Interactive serving has far less time per request, so redirecting work needs spare capacity and recoverable request state, and substituting a smaller model needs an acceptable quality fallback. SLOs define which failures and delays the service may tolerate, and a recovery mechanism helps only when its added work fits within that contract.

The same targets also define the serving system's capacity metric. **SLO goodput** is the rate of completed requests that meet the stated TTFT and TPOT targets for a fixed request mix. It is the serving counterpart of the useful-progress goodput that coordination protects in training (section 1.5.1.3). ML goodput measures how much of the fleet's capacity becomes useful work at the ideal rate; SLO goodput counts only the part of that capacity that completes requests within their targets. A request that misses its contract counts for nothing, however busy the accelerators were while serving it, so increasing arrivals without completing more requests within those targets does not increase useful capacity.

10.3.2 The serving tax

↩ ML goodput, the fraction of fleet capacity that becomes useful work at the ideal rate, and its three factors are established in section 9.2.

If an interactive request carries a 100 ms latency target and model execution takes 70 ms on an idle replica, routing, network transfers, serialization, coordination, and queueing must all fit inside the remaining 30 ms. Distributing inference adds transfers and coordination to a request that already spends time in host processing and queues. This "serving tax" is the exposed work outside model execution that consumes the request's latency budget, and its size depends on which components the request crosses and which costs can overlap.

Network communication contributes setup and transfer time at each cross-machine dependency. Synchronization between GPUs can expose that cost when the next operation needs a remote result. For an illustrative path with four sequential exchanges taking 50–500 microseconds each, the exposed contribution is 200 microseconds to 2 milliseconds. Actual collective time also depends on payload,

topology, contention, and the number of communication rounds; a generic network round-trip time cannot be charged once per collective without checking its implementation.

↩ The α-β model, which splits each transfer into a fixed startup latency plus a bandwidth-dependent term, is established in section 3.4.4.

Before and after those network hops, serialization adds host work whenever a service converts structured data into a wire representation and reconstructs it at the receiver. Protocol Buffers, JSON, and Python's `pickle` provide different representations and implementation paths, so their cost must be measured for the actual payload and runtime. Tensor transfers within a sharded model need not pass through the same object-serialization path as a service request. Formats such as FlatBuffers and Cap'n Proto permit access to wire-format data in place, avoiding a separate unpacking step when the consumer supports that representation, though they do not eliminate network transfer or every downstream copy.

Once the payload is on the wire, load balancing adds a routing decision and sometimes another forwarding hop: examining request metadata, consulting replica state, and forwarding to the selected backend all consume time before execution begins. The measurement must separate decision time from waiting and transport so that a slow routing path is not mistaken for a slow model.

When a request fans out to several backend services such as feature stores, retrieval indexes, and ranking models, coordination overhead emerges as the gateway joins their results. Independent feature or model calls can run concurrently, but a dependent ranking stage must wait for the inputs it needs, so the critical path includes exposed coordination and the slowest required branch rather than the sum of every branch's runtime. Finally, queuing delay ($L_{\text{lat,queue}}$) adds waiting time whenever work arrives while the required resources are busy, even when the average arrival rate remains below average service capacity.

The serving tax must be measured on the same request boundary as model execution. Equation 10.1 decomposes that latency, with each term counting an exposed contribution so concurrent work is not counted twice:

$$L_{\text{lat,total}} = L_{\text{lat,infer}} + L_{\text{lat,network}} + L_{\text{lat,serialization}} + L_{\text{lat,coordination}} + L_{\text{lat,queue}} \tag{10.1}$$

In the 100 ms request budget above, the four sequential exchanges alone can claim up to 2 ms of the 30 ms non-compute slack before any request waits for a busy replica. The measured contribution determines the intervention: co-location removes a network hop, fewer dependent exchanges reduce startup costs, and a faster link shortens large transfers, whereas queueing requires controlling admitted work or increasing usable capacity. Each choice spends resources to recover latency budget, connecting the request-level diagnosis to lifetime serving cost.

10.4 Serving Cost Dominance in Production

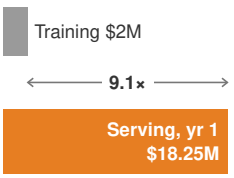

At this traffic, one year of serving costs about nine times the training run.

A team that budgets millions of dollars to train a model can find within weeks of launch that serving traffic is spending more compute per month than the entire training run cost. While the serving tax in equation 10.1 consumes a fraction of each request's latency budget, those per-request costs compound across every query over the model's operational lifetime. By the serving cost dominance principle (17), recurring serving expense can dwarf upfront training expense by orders of magnitude because a model version incurs training cost once before deployment while serving cost recurs with traffic. Both activities can consume purchased or rented infrastructure; the distinction here is expense timing, not financial capitalization. In a high-volume deployment, that recurring traffic bill rapidly overtakes the one-time training investment (notebook 10.1).

> **Napkin Math 10.1: The serving cost multiplier**
>
> **Problem**: A team has spent $2,000,000 training a 70B parameter model and now serves it to 1,000,000 daily active users, each making 50 requests/day. Is training or serving the dominant cost over 1 year?
>
> **Math**:
>
> 1. **Training cost**: $2,000,000 (one-time).
> 2. **Metric**: Annual inference volume is 10^6 users $\times$ 50 reqs $\times$ 365 days $\approx$ 18.25B requests/year.

3. **Cost per request**: Assume a 70B model on H100 costs ≈ \$0.001/request (input + output).
4. **Annual serving cost**: 18.25B requests × \$0.001/request = \$18,250,000.

Systems insight: Serving costs 9.1× more than training in the first year alone. A 10 percent reduction in cost per request, at unchanged request volume and service quality, saves about \$1.8M in the first year, nearly paying for the original training run.

For one model version at a fixed query rate and average cost per served query, equation 10.2 combines training expense with recurring serving expense:

$$C_{\text{total}} = C_{\text{train}} + C_{\text{serving}} \times T_{\text{deployment}} \times \lambda_{\text{arr}} \tag{10.2}$$

where C_{train} is the one-time cost to train the model, C_{serving} is the cost per query served, $T_{\text{deployment}}$ is the deployment duration in appropriate time units, and λ_{arr} is the query rate. The rate must count the same served queries used to estimate unit cost. That cost must include the chosen serving boundary, including idle or reserved capacity when it is billed. Changing traffic mix, fleet size, or price requires evaluating separate intervals rather than treating one unit cost as constant throughout deployment.

The same leverage can hold across application categories with different cost structures. In the chapter's illustrative recommendation scenario, a DLRM trained for \$8,000 (\$2,000 hardware for 1,000 GPU-hours at \$2/GPU-hour, plus 3× additional engineering cost for data and experimentation) and served at 10,000 QPS for 730 days accumulates 6.31×10^{11} lifetime queries. At the assumed \$10/million queries, those queries cost \$6,307,200 to serve, producing a 788.4× ratio in this scenario. Other workload envelopes span orders of magnitude in both terms, and the sensitivity ranges in table 10.3 are illustrative envelopes rather than market-price estimates.

Table 10.3: Illustrative Training and Serving Cost Ranges: Independent sensitivity ranges for training a model version and serving it for a year. A cost ratio requires a paired training-and-serving scenario; it cannot be inferred from matching endpoints of independent ranges. A deployment must substitute measured traffic and a consistent accounting boundary.

Application	Training cost	Annual serving cost
Recommendation (high QPS)	\$10K–\$100K	\$1M–\$10M
Search ranking	\$100K–\$1M	\$10M–\$100M
LLM API	\$1M–\$100M	\$10M–\$1B
Internal analytics	\$1K–\$10K	\$10K–\$100K

A reduction in cost per completed request yields recurring savings only if workload volume and service quality stay fixed; lower latency alone does not reduce billed capacity. Unit cost also changes with daily request volume, as the constructed cost curves of figure 10.2 illustrate. Against its assumed price of \$0.05 per 1,000 tokens, the optimized curve crosses near 2×10^5 requests per day and the unbatched single-GPU curve near 4×10^7 requests. These crossings follow from the assumed curves, not measured profitability, so a deployment needs its own token mix, occupied capacity, and cost boundary to locate its break-even point.

The cost ranges in table 10.3 give a static comparison. Figure 10.3 instead accumulates assumed monthly serving expense over a 36-month deployment and marks when the serving portion equals the initial training expense, within days in the plotted recommendation scenario and roughly six weeks in the 70B LLM scenario. These timings follow from the specified expenses, not from model size or user count, and a larger training expense, lower traffic, or shorter deployment can reverse which activity dominates. When recurring traffic makes inference the dominant expense, serving-specific designs can substantially reduce cost and power per query (Patel et al. 2024), provided each optimization targets the workload's actual bottleneck and the right level of the serving stack.

10.5 Serving Workloads and the Serving Hierarchy

An optimization tuned for one model family or one layer of the stack can miss the active bottleneck when applied to another. Paging a key-value cache does nothing for an embedding-heavy recom-

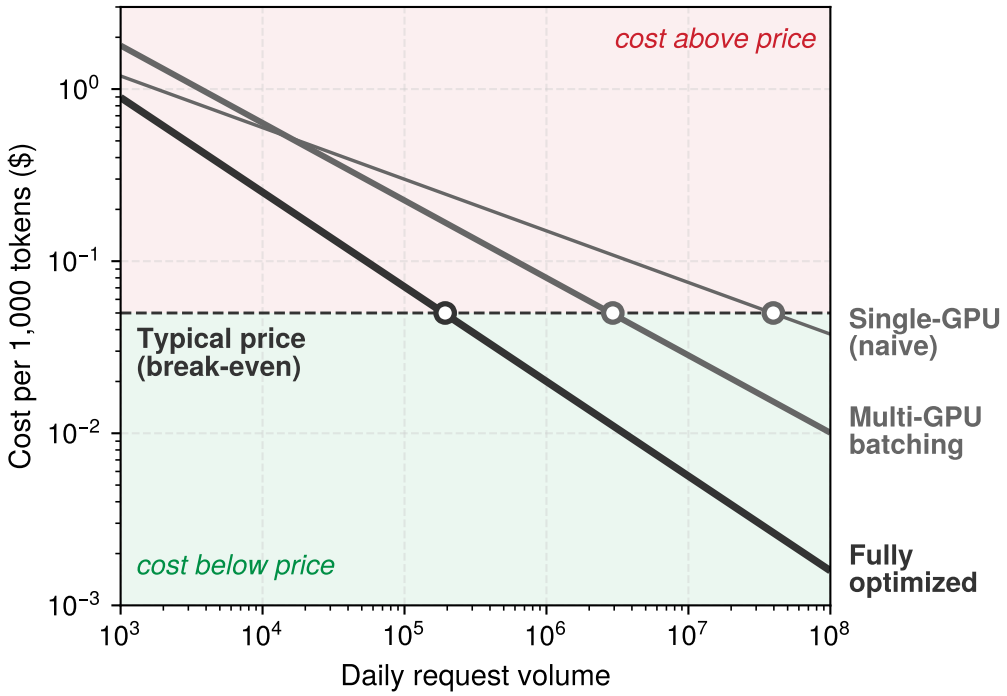

Figure 10.2: The Serving Cost Curve: Constructed unit-cost curves cross an assumed price line at different request volumes. The optimized curve crosses earlier under these assumptions; the crossings do not establish measured profitability or guaranteed savings from quantization and batching.

mendation model, and fusing a GPU kernel cannot rescue a service whose requests stall in a load balancer. Controlling serving cost therefore requires separating the workload mix from the system layer that owns each bottleneck.

10.5.1 Inference beyond LLMs

Designing a serving platform around autoregressive token generation alone leaves a fleet unprepared for the workloads that dominate consumer request volume. In large consumer platforms, recommendation and ranking workloads can consume the majority of AI inference cycles and capacity (Gupta et al. 2020; Hazelwood et al. 2018), stalling on terabyte-scale sparse embedding lookups under sub-10 ms deadlines rather than sequential weight reads. Vision and image processing, NLP and LLM workloads, fraud detection, ads, and other classification tasks create additional serving demands. Table 10.4 compares representative pressures: sparse lookup for recommendation, batch efficiency for fixed-shape vision, weight and cache traffic for LLM decode, and incremental deadlines for speech. These are starting points for diagnosis; the active bottleneck still depends on the model, input shape, and deployment.

Request volume must be derived from application behavior. In a recommendation service, a page load, scroll, or click can trigger ranking or retrieval, so one browsing session can produce many inference requests. A language-model application may likewise issue several model calls for one user action or run background tasks without an explicit query. Capacity planning therefore counts the model work triggered by each interaction rather than treating user count as request rate.

Workload structure also determines which techniques can help. Embedding sharding and lookup locality address large recommendation models (Naumov et al. 2019; Gupta et al. 2020), and feature store architectures must deliver the inputs a model expects within its latency budget, while continuous batching and KV cache management address variable-length autoregressive work. Text-to-image systems such as DALL-E add multimodal serving demands (Ramesh et al. 2021) whose volume and latency targets remain product-specific.

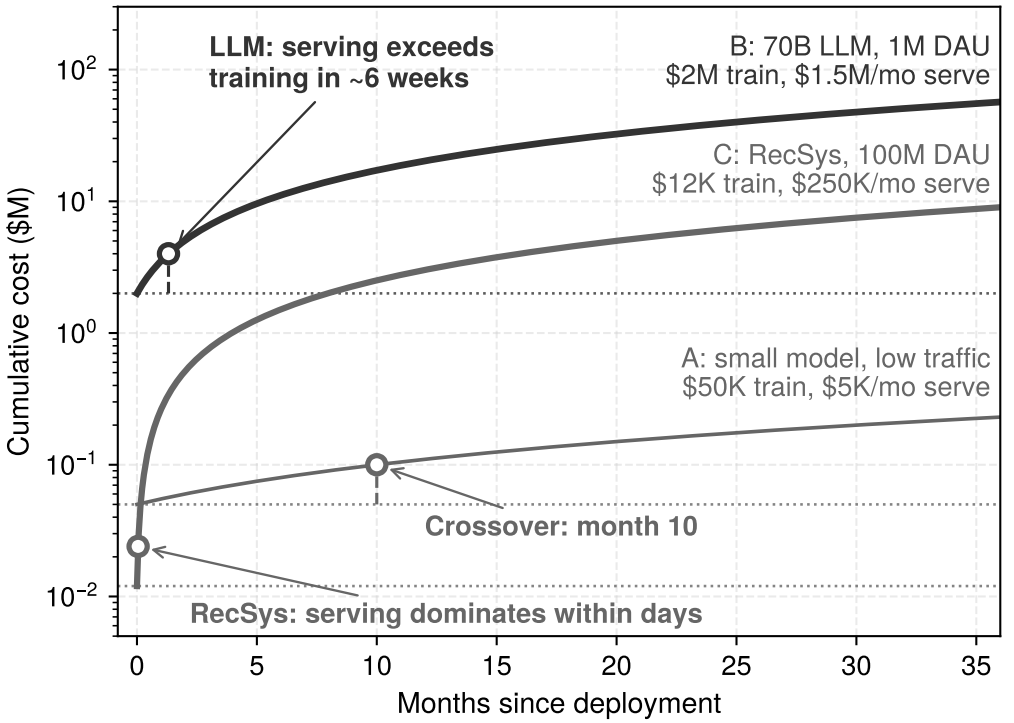

Figure 10.3: The Serving Cost Crossover: Cumulative total deployment cost (initial training plus ongoing serving) for three illustrative scenarios. Vertical markers (scenarios A and B) indicate when the serving portion equals the one-time training cost. In the two high-traffic scenarios, serving expense matches training expense within weeks or days. In the low-traffic scenario, training expense remains larger over the plotted period. These differences follow from the specified training costs, traffic, and unit serving prices.

Table 10.4: Production Inference Landscape: Different model types create different serving pressure, latency requirements, and optimization challenges. The table is a qualitative workload map rather than a universal traffic-share distribution. Technique selection must match the specific workload.

Model type	Serving pressure	Latency target	Key challenge
Recommendation/ranking	High volume in consumer platforms	$<$ 10 ms p99	Embedding lookup
Vision (CNN)	Workload-dependent	20–100 ms	Batch efficiency
LLM	High cost per request	100 ms–10s	Memory bandwidth
Speech/Audio	Stream-driven	Real-time	Sequential decode
Multimodal/text-to-image	Emerging and product-specific	Varies	Cross-modal coordination

10.5.2 The serving hierarchy

Tuning a replica's kernels or cache allocator cannot fix tail latency when requests are waiting inside a router or starving behind a noisy neighbor on a shared cluster. To locate the right control, serving optimizations organize into a four-level serving hierarchy ordered by scope, where each level owns a distinct bottleneck and measurement boundary. The request level follows one request through preprocessing, batch-formation waiting, response caching, and model execution, where the target is the latency a user sees. The replica level looks inside one model instance, where utilization, KV-cache memory management, and kernel efficiency set the throughput it can deliver before it saturates. The service level turns many replicas of the same model into aggregate capacity through load balancing, request routing, and autoscaling. The platform level decides, across services and tenants, whether a shared serving fleet stays efficient without allowing one workload to degrade another.

Each level supplies a different measurement boundary. Table 10.5 associates request latency, replica throughput, aggregate service capacity, and cross-tenant efficiency with the controls most directly responsible for them. Their effects still cross boundaries: batching can increase replica throughput while adding request delay, and replica placement can change both communication cost and service capacity. The hierarchy helps locate the decision without treating its metric in isolation. In C^3 terms, batching and cache management improve compute reuse, sharding introduces communication, and routing, isolation, and autoscaling coordinate the resources that make those gains usable.

Table 10.5: Serving Hierarchy Optimization Targets: Each level of the serving hierarchy targets distinct performance metrics, from per-request latency at the request tier to aggregate resource efficiency at the platform tier.

Level	Optimization target	Key techniques
Request	Per-request latency	Dynamic batching, response caching
Replica	Throughput, utilization	KV-cache paging, kernel fusion
Service	Aggregate capacity	Load balancing, routing, autoscaling
Platform	Resource efficiency	Multi-tenancy, scheduling, placement

Checkpoint 10.2: The serving hierarchy

Verify your understanding of where specific optimizations sit within the serving hierarchy:

- ☐ **Replica level**: identify the hierarchy level at which KV-cache paging reduces GPU memory waste, and explain why that raises concurrent request capacity rather than single-request latency.
- ☐ **Caching level**: distinguish response caching, which shortens one request's path, from KV-cache paging, which raises a replica's capacity, and place each in the hierarchy.
- ☐ **Load balancing**: identify why load balancing belongs at the service level and explain how its choice of replica can affect both aggregate goodput and a request's queueing delay.
- ☐ **Isolation level**: evaluate the claim that multi-tenancy and resource isolation are primarily request-level concerns.

The serving decisions that follow climb the hierarchy from the request level to the platform level. Batching sets how long each request waits to share execution. KV-cache layout, decode-time policies, and adapter state decide how many requests a replica can hold. The split between prefill and decode decides whether a new prompt and the streams already decoding share a worker or separate pools. Weight quantization and model sharding change the replica's memory and compute budget. Routing, isolation, autoscaling, and global placement then turn replicas into a service and a platform. Precision enters every capacity calculation that follows through two separate budgets: reducing weight precision frees resident memory, whereas reducing KV precision lowers the memory required per active sequence. Section 10.9 compares the serving trade-offs of these representation choices before sharding divides a replica across devices.

10.6 Batching Strategies at Scale

A stream of hundreds of disparate chat requests arrives at an inference server every second. Processing them one by one leaves much of the GPU idle, starved for work between small decode steps. Waiting to gather a large static batch forces the first request in line to endure unacceptable latency. Batching at scale therefore needs policies that group requests as they arrive without violating tail-latency deadlines.

Processing multiple requests together shares weight reads and kernel launch overhead across more useful work. The resulting service-rate gain can reduce queueing, but waiting to form a batch

adds delay; the net effect on request latency depends on load. Single-machine serving applies this insight through dynamic batching, which collects requests within a time window before processing them together.

At scale, batching must also account for workload structure. A policy suited to uniform vision inputs can leave slots unused during variable-length generation, while grouping sparse embedding lookups solves a different problem from batching dense execution. The unit of reusable work determines which policy is worth evaluating. The strategies below proceed in order of mechanism rather than traffic share: vision (request batching), LLMs (iteration-level batching with KV state), and recommendation (feature-parallel batching over distributed embeddings), even though recommendation often dominates consumer inference volume (section 10.5.1).

10.6.1 Batching across model architectures

Holding arriving requests until a fixed batch fills forces early arrivals to wait and pads shorter inputs with wasted work, whereas running every request immediately leaves later arrivals queued behind one-request passes that underuse the GPU. Batching efficiency depends on how computation scales with batch size relative to memory and communication constraints, and the execution timeline in figure 10.4 separates waiting to form a batch from waiting for the accelerator to become available. Unbatched serving adds no batch-formation delay, but later requests can still queue behind earlier forward passes. Static batching consolidates all four illustrated requests into one pass, adding formation delay and padding to accommodate different input lengths. Dynamic batching closes formation windows at fixed time intervals (Δt), producing two passes with two requests each in this example. The window limits intentional formation delay; it does not bound total queueing when ready batches arrive faster than the server can execute them.

In vision models processing fixed-size images, arithmetic work grows roughly in proportion to batch size while requests share the weights, so larger batches improve weight reuse until compute, memory, or the latency budget binds. The largest feasible batch is not necessarily the best interactive operating point. By contrast, in LLMs during the decode phase, small batches perform little computation per byte of active weights loaded, so memory bandwidth binds. Larger batches raise arithmetic intensity until compute, KV-cache traffic, or the latency budget limits further gains. Small-batch decode therefore does not inherit a new accelerator generation's headline compute improvement. From Volta to Blackwell, peak compute climbs 18× while memory bandwidth expands by only 8.9×, and figure 10.5 shows generation throughput tracking the slower curve whenever weight streaming dominates at small batch sizes.

Recommendation systems shift the bottleneck again: sparse embedding lookup can limit throughput before dense computation does. Grouping sparse accesses improves locality and distributes lookup work across shards, while batching the dense stages improves matrix multiplication throughput. The serving policy must therefore account for both stages and the communication between them rather than infer the bottleneck from the model category alone.

↩ The accelerator trajectory behind this divergence, with compute outpacing memory bandwidth from one generation to the next, is established in section 2.2.3.

10.6.2 Batching physics and efficiency curves

Raising batch size to maximize a replica's throughput stretches each batch's execution time and holds arriving requests in a formation queue until the service violates its latency SLO. Batch size is the single-replica control knob that moves a workload along the roofline by reusing each weight read across more requests, which raises arithmetic intensity. At fleet scale, the same knob also sets how long arriving work waits and how much request state the replica must hold. A model relating batch size (B), batch service time ($T_{\text{svc}}(B)$), and batched service rate ($\mu_{\text{eff}}(B)$) exposes that trade-off; selecting an operating point also requires the arrival pattern, memory budget, and latency target.

The first model isolates batch execution. Here $T_{\text{svc}}(B)$ is the service time once a batch begins execution; it excludes queueing and batch formation. Keeping the measurement boundary explicit prevents an execution-time improvement from being mistaken for an end-to-end SLO guarantee.

↪ The roofline bound, and why batch-one decode sits far below the ridge point, are collected for reference in section A.3.1.

The service-time equation decomposes batch execution time into fixed overheads (kernel launch, memory loading) and variable costs (compute per sample):

$$T_{\text{svc}}(B) = T_{\text{fixed}} + B \times t_{\text{req}}$$

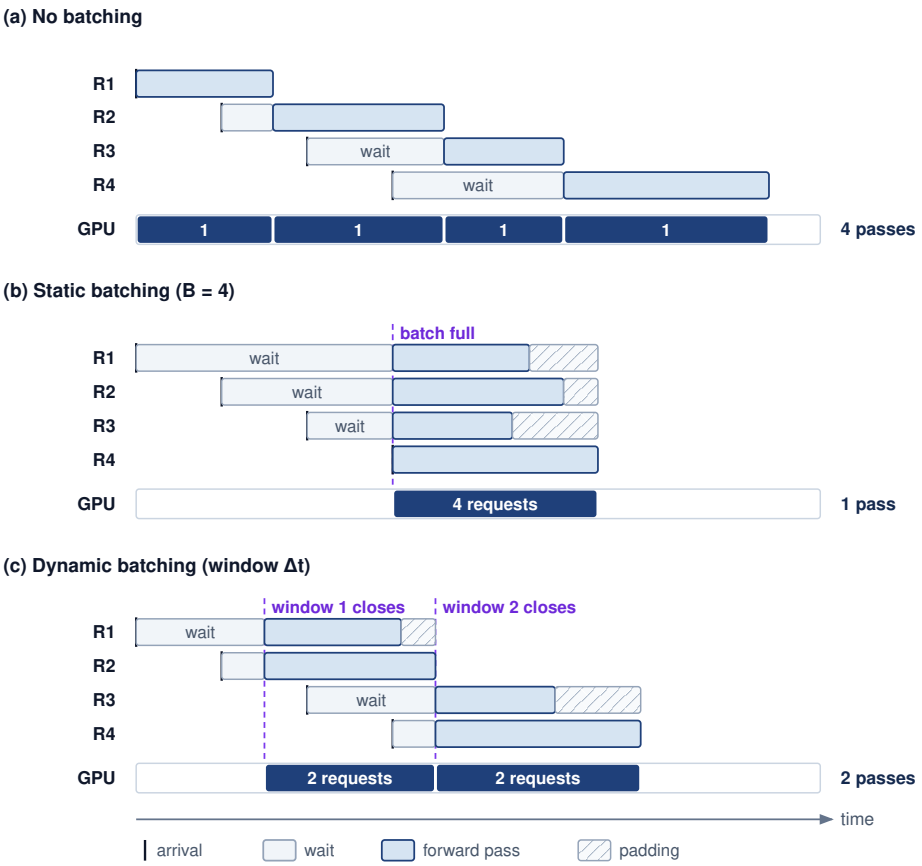

Figure 10.4: Batching Strategies Compared: Four requests of different lengths arrive one time unit apart under three policies; the timings are illustrative. (a) Without batching, each request receives its own forward pass, so later arrivals queue behind earlier ones and the GPU runs four passes. (b) Static batching holds every request until the batch of four is full, then runs one pass padded to the longest request; early arrivals pay the longest wait. (c) Dynamic batching dispatches whatever has arrived when a short window closes, limiting intentional formation delay in this timeline while sharing each pass between two requests. Waiting for an occupied worker remains a separate source of delay.

In this formulation, T_{fixed} captures costs paid once per batch, such as loading weights from high-bandwidth memory (HBM) and kernel launch latency, while t_{req} denotes the marginal cost of adding one request, such as the compute time for that sample. Treating these terms as constants is a local approximation for a fixed workload and execution regime. Kernel changes, cache traffic, and resource saturation can alter the measured curve as batch size grows.

The throughput equation describes the system's capacity:

$$\mu_{\text{eff}}(B) = \frac{B}{T_{\text{svc}}(B)} = \frac{B}{T_{\text{fixed}} + B \times t_{\text{req}}}$$

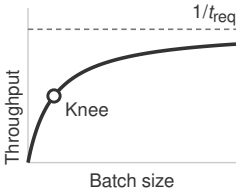

Throughput saturates toward $1/t_{req}$, past the knee, each larger batch buys less.

The resulting batching efficiency curve shows diminishing returns from amortizing overhead. For small batch sizes (B), T_{fixed} dominates service time, so increasing B shares that cost and raises throughput substantially. As B grows, the relative contribution of T_{fixed} shrinks and the model approaches its asymptotic rate $1/t_{\text{req}}$. This limit alone does not identify whether compute or bandwidth binds on the actual device. The knee of the curve, where fixed and variable costs become comparable, is a useful candidate operating point. It is not an optimum independent of the service contract; a tight latency budget may require a smaller batch, while a throughput-oriented workload may benefit from a larger one.

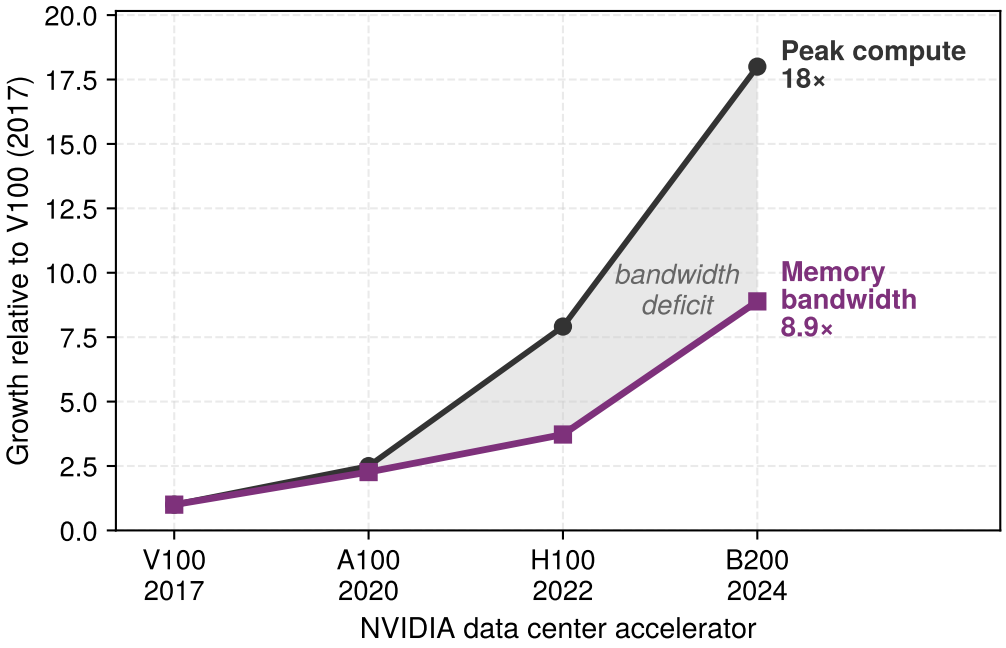

Figure 10.5: Decode Inherits the Bandwidth Deficit: Peak compute and memory bandwidth for four accelerator generations, normalized to the V100. Compute rises 18× against bandwidth's 8.9×. In weight-bandwidth-bound decode, each iteration loads the active weights and reuses them across its output tokens; at fixed batch size, bandwidth growth constrains the gain. These hardware ratios are not measured workload speedups. Specifications are dense peak rates and bandwidths from the vendor datasheets recorded in the chapter's data file.

The execution-only feasibility condition is $T_{\text{svc}}(B) \leq \text{SLO}$. In a live service, execution receives only the budget remaining after routing, queueing, batch formation, and response delivery. The batch-size decision must therefore use that remaining budget rather than consume the full SLO. Comparing a vision workload with a relatively high t_{req} against an LLM decode workload with a high T_{fixed} from weight loading shows how those measured costs shift the efficiency knee (notebook 10.2).

Napkin Math 10.2: The batching efficiency curve

Problem: An engineer optimizes two services: a vision model (ResNet) and an LLM (70B). At what batch size does each hit the "knee" of its efficiency curve?

Math: The knee occurs where the variable compute cost $(B \times t_{\text{req}})$ starts to exceed the fixed overhead (T_{fixed}).

1. **Vision Model** $(T_{\text{fixed}} = 2\text{ ms}, t_{\text{req}} = 1\text{ ms})$:
   - Knee: $B \approx 2\text{ ms}/1\text{ ms} = 2$.
   - **Result**: This simplified latency model reaches its first overhead-amortization knee at a small batch. It does not locate hardware saturation; measuring throughput at larger batches tests whether additional concurrency still helps.
2. **LLM Decode** $(T_{\text{fixed}} = 40\text{ ms}, t_{\text{req}} = 0.5\text{ ms})$:
   - Knee: $B \approx 40\text{ ms}/0.5\text{ ms} = 80$.
   - **Result**: The larger fixed cost in this decode example moves the overhead-amortization knee to a larger batch.

- **Where 0.5 ms comes from**: It is an assumed marginal cost per added sequence, dominated by reading that sequence's own KV state at a moderate context length. Unlike the weight stream, that read is private to each sequence, so it grows with batch size instead of being shared.
- **Where 40 ms comes from**: It is roughly the time to stream the model's 140 GB of FP16 weights through one H100's 3.35 TB/s of HBM bandwidth, the aggregate weight traffic of one decode step expressed in single-device time. No single H100 holds those weights. The reference replica of section 9.1 spreads them over eight GPUs, which divides both T_{fixed} and t_{req} by the tensor-parallel degree and leaves their ratio, and therefore the knee, unchanged.

Systems insight: Because streaming the weights is the dominant per-step cost of decode, the system has not reached the efficiency knee until batch size 80. The knee depends on the ratio of fixed to per-request cost, not on how many devices share them, which is why sharing weight reads across concurrent requests matters at every replica size. The batch a service actually runs also depends on context length, cache capacity, kernel efficiency, and the latency target; continuous batching and paged allocation are the mechanisms that let a replica reach it.

While the analytical batching model locates the knee from each model class's ratio of fixed to per-request cost (table 10.9), production serving stacks measure the resulting throughput-latency frontier directly. The open InferenceX benchmark serves DeepSeek-R1, a mixture-of-experts reasoning model, with 1,024-token prompts and 1,024-token responses on each accelerator across many combinations of serving framework, parallelism, and concurrency (SemiAnalysis 2026). For every configuration it records the output tokens per second each user receives, called interactivity (the reciprocal of TPOT), and the output tokens per second each GPU delivers across all users. Figure 10.6 plots 325 such configurations at FP8 precision and traces, for each accelerator, the frontier of configurations that no other configuration beats on both rates. Every frontier falls as interactivity rises, because a faster stream per user permits fewer concurrent sequences per decode step. The best B200 configuration delivers about 6,300 output tokens per second per GPU when each user receives about 13 tokens per second. It delivers about 210 when each user must receive at least 100, a 30× difference in per-GPU output paid for faster streams. The interactivity target alone moves per-GPU output by that much, so it must be fixed before any two throughput figures can be compared.

The same frontiers separate hardware generations by more than their datasheets do. At 50 tokens per second per user, the best measured H100 configuration delivers about 80 output tokens per second per GPU and the best GB200 NVL72 configuration about 3,650, a 44× difference. Yet the peak dense FP8 rate of a GB200 GPU is only about 2.3× that of an H100. The benchmark does not apportion the gap, but the candidates lie mostly beyond peak arithmetic, in memory bandwidth and capacity per GPU, a 72-GPU NVLink domain that keeps experts and KV state off slower networks, and platform-tuned serving software. A throughput figure quoted without its latency target therefore says little about either system.

Translating the batch-level equations into system-wide capacity planning requires Little's law,[4] a classical result from queuing theory that relates concurrency, throughput, and latency in any stable system. Whereas $T_{\text{svc}}(B)$ prices one replica's execution batch, Little's law connects the external arrival rate and mean request time in system to the total number of requests concurrently in flight across the fleet.

[4] **Little's law**: The identity holds for any stable system regardless of the arrival or service distribution, which is why it survives the bursty, variable-length traffic of inference. Because it relates averages, it bounds neither maximum service capacity nor tail latency.

↩ The M/G/1 waiting-time model, in which utilization and service-time variability multiply queueing delay, is established in section 8.3.

Theorem 10.1: Little's law for inference

Concept: In any stable queuing system, the average number of requests in the system (Q_{req}) equals the arrival rate (λ_{arr}) multiplied by the average time a request spends in the system (T_{lat}).

$$Q_{\text{req}} = \lambda_{\text{arr}} \cdot T_{\text{lat}}$$

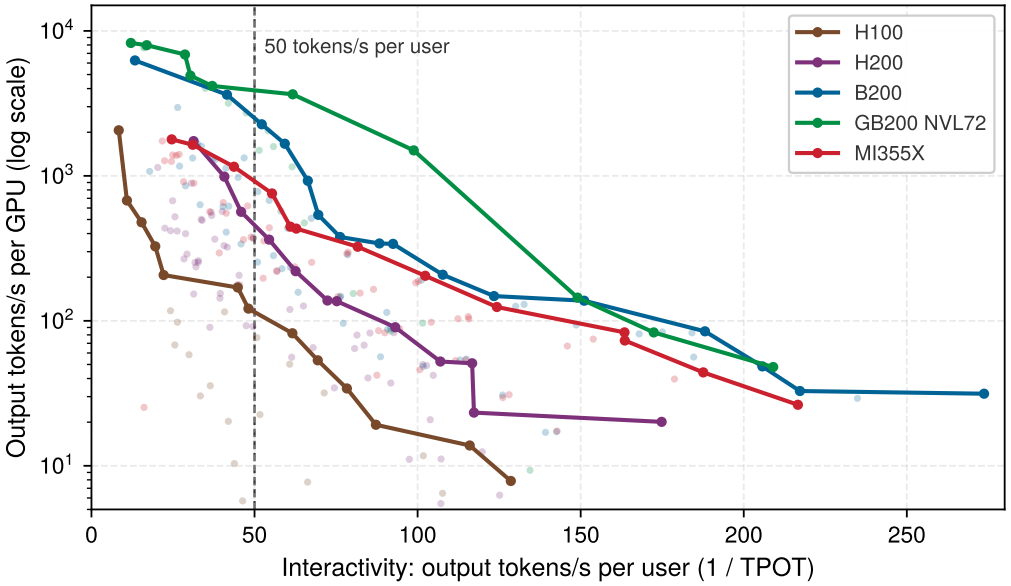

Figure 10.6: Throughput versus Interactivity in Measured LLM Serving: Output tokens per second per GPU against output tokens per second per user for DeepSeek-R1 at FP8 with 1,024-token prompts and responses. Faint points are individual serving configurations (framework, parallelism, and concurrency); each line is the frontier of configurations that no other configuration on the same accelerator beats on both axes. The dashed line marks 50 tokens per second per user. Data from SemiAnalysis (2026), latest published result per configuration as of October 2026.

Application: Concurrency planning. This idealized service counts only batch formation, queueing, and execution in its end-to-end latency; routing and response delivery add no delay. Assume the mean latency given below and full execution batches of the stated size. A percentile SLO cannot generally be substituted for the mean in Little's law.

- **Target Throughput** (λ_{arr}): 1,000 requests/sec
- **Assumed mean latency** (T_{lat}): 100 ms (0.1 s)

Mean in-flight work (Q_{req}): $1000 \times 0.1 = 100$ concurrent requests

Capacity planning: If a single GPU replica executes batches of 8 in 80 ms, and the fleet targets at most 80 percent mean service utilization:

1. Replica throughput $= 8/0.08 = 100$ req/s
2. Throughput-only lower bound $= \lceil 1000/100 \rceil = 10$ replicas
3. Illustrative headroom-sized replica count $= \lceil 1000/(0.8 \times 100) \rceil = 13$ replicas

Verification: Maximum executing positions $= 13$ replicas $\times\ 8$ batch positions $= 104$. These positions do not include requests waiting for execution or batch formation. Under the assumed fixed batch service time, Little's law on the execution boundary gives 80 requests executing on average; the remaining 20 of the assumed mean in-flight population are waiting or accumulating. The throughput-only lower bound would run at the stability boundary; the utilization-sized candidate runs at about 76.9 percent service utilization. The assumed mean latency still needs validation, and this arithmetic does not prove a percentile latency SLO; that requires an arrival/service model or a measured load test.

A load test can locate the stability boundary, where arrivals equal service capacity, only when its load generator keeps offering work at the stated arrival rate. The measurement contract of section 9.3.1 guards against load that eases off while the replica stalls. A serving measurement therefore holds a fixed external arrival process rather than a fixed concurrency, because response-

paced generation can reduce arrivals during a stall and make the system appear to have more headroom than it does.

↪ The mean queuing delay under a batched M/G/1 surrogate, and how utilization headroom and batch size trade against latency, is developed in section F.1.1.

10.6.3 Static and dynamic batching for vision models

Static batching, which waits until exactly B requests arrive before launching a forward pass, leaves early requests stranded during traffic lulls until their latency budgets expire. Vision models otherwise present the simplest batching case: preprocessed inputs share a fixed tensor shape and execute in a predictable time, so single-machine batching applies directly and fleet scale adds only the question of where batches form.

Dynamic batching collects requests for a maximum time window T_{window} or until reaching maximum batch size B_{max}, whichever occurs first. For Poisson arrivals with rate λ_{arr} (requests arriving independently of one another at that constant average rate), equation 10.3 organizes expected request latency into queueing, formation, and execution contributions:

$$\mathbb{E}[T_{\text{lat}}] = \mathbb{E}[L_{\text{lat,queue}}] + L_{\text{lat,form}} + T_{\text{svc}}(B) \tag{10.3}$$

where $\mathbb{E}[L_{\text{lat,queue}}]$ is the expected queuing delay, $L_{\text{lat,form}}$ represents mean batch formation delay (up to T_{window}), and $T_{\text{svc}}(B)$ represents mean inference time over the realized batch sizes B. Queueing here excludes time already counted as batch formation. The two waiting terms respond differently to load. Under Poisson arrivals the mean inter-arrival gap is $1/\lambda_{\text{arr}}$, so a higher arrival rate fills a batch faster and reduces formation delay. Queueing for a busy replica can nevertheless grow as arrival rate approaches service capacity; the batching timeout does not cap that queueing delay. Spending a small slice of the latency budget on batch formation can multiply per-GPU capacity while keeping end-to-end latency within target (example 10.1).

> Example 10.1: Dynamic batching for ResNet-50
>
> **Scenario**: A vision classification service receives requests at 5,000 QPS with a 50 ms p99 latency SLO across 10 replicas, so each replica sees 500 QPS. One image served alone takes 5 ms. The scenario assumes that a batch of about 5 images executes in 8 ms and a batch of about 10 in 12 ms, because a batch shares weight reads and kernel launches across its images.
>
> **Diagnosis**: Each utilization below is the per-replica arrival rate multiplied by the service time one request consumes. Option A serves every image alone, which overloads the replica ($\rho_{\text{serv}} = \lambda_{\text{arr}} \times T_{\text{svc}} = 500 \times 0.005 = 2.5$ utilization). Option B holds a 10 ms batching window, which collects about 5 requests (arrival rate $\times$ window). Their batch runs in 8 ms, so each request consumes 1.6 ms of execution, and utilization falls to $\rho_{\text{serv}} = 500 \times 0.0016 = 0.8$. The wider 20 ms window of Option C collects about 10 requests whose batch runs in 12 ms, or 1.2 ms each, reaching $\rho_{\text{serv}} = 500 \times 0.0012 = 0.6$ and a further 33 percent of per-replica capacity. Both options leave service headroom. Whether they also meet the p99 SLO depends on the window plus the queueing delay each utilization implies under the service's real arrival pattern, which utilization alone does not determine.
>
> **Systems lesson**: Dynamic batching is effective when the latency budget has unused slack. The serving policy converts latency headroom into throughput, but requires careful tuning to prevent SLO violations.

↪ A procedure for this tuning, from fitting the service-time model to checking the p99 target, is developed in section F.1.1.5.

At scale with multiple replicas, batch formation can occur either at individual replicas or at a centralized batching layer. Replica-local batching has each replica independently form batches from its assigned traffic. This avoids a shared batching stage, but splitting a light arrival stream among many replicas can leave each waiting to fill its batch. Centralized batching pools arrivals before dispatching formed batches, potentially shortening formation waits, at the cost of an additional network hop and a shared scheduling bottleneck. Load balancing helps local batching, but equal request counts do not guarantee equal work or reproduce the formation benefit of pooling arrivals. The choice depends on arrival density, request compatibility, and replica service times.

10.6.4 Continuous batching for LLM inference

When a request-level batch locks its membership until the longest sequence finishes, a short reply that ends after 10 tokens leaves its slot stranded for the remaining 90 steps of a 100-token neighbor. Queued requests starve while the GPU still streams the full model weights on every step. The autoregressive bottleneck (principle 18) makes that wasted slot expensive: at small batch sizes, decode is typically bandwidth-bound because each iteration streams active model weights while generating only one token per sequence. Batching shares those weight reads across concurrent requests, and as the batch grows, compute or KV-cache traffic can become the binding constraint, so bandwidth-bound decode is an operating regime rather than a property of every serving configuration.

Definition 10.1: Continuous batching

Continuous batching is a serving strategy that updates batch membership at iteration boundaries, allowing new requests to enter without waiting for every request in the original batch to finish.

1. **Significance**: It reduces padding waste by refilling completed slots at iteration boundaries. Throughput gains require waiting requests, available KV-cache capacity, and scheduling overhead small enough to preserve the benefit.
2. **Distinction**: Unlike static or dynamic batching (which group requests at the request level), continuous batching operates at the iteration level, reshaping the active batch between decode iterations.
3. **Common pitfall**: Continuous admission must coordinate scheduling with KV-state ownership and lifetime. A paged KV-cache allocator such as PagedAttention (section 10.7.3) improves space utilization, but paging is not required: an implementation can recycle fixed-capacity request slots, at the cost of unused reserved space.

The Orca system (Yu et al. 2022) removed the request-level barrier by updating membership at iteration boundaries to enable slot reuse. Once a request completes, another can enter if its prompt has been processed and sufficient cache capacity is available. The scheduler must account for that preparation rather than treat every replacement as a decode-only operation.

Consider a fixed-width batch of 8 sequences. If one sequence completes after 10 tokens while others require 100 tokens, its slot contributes no useful output for 90 iterations. If the kernel continues executing padded work for that slot, the fraction of wasted compute attributable to it is:

$$\text{Wasted compute} = \frac{(100-10)\times 1}{100\times 8} = 11.25\%$$

The waste grows with the spread of output lengths, a dependence that section 10.6.5 turns into a formula. A kernel that skips completed sequences avoids their padded arithmetic, but a fixed-membership policy still prevents waiting requests from using the vacated capacity. Algorithm 10.1 states the continuous-batching scheduler as a decode loop with a prefilled waiting queue. In this simplified policy, completed requests release their KV-cache pages (the fixed-size blocks in which cache memory is allocated), eligible waiting requests enter freed slots, and the next kernel runs over the updated active set. For Archetype A workloads (GPT-4/Llama-3), refilling slots lets more useful tokens share each weight read. The gain comes from reusing transferred weights, not merely from saturating bandwidth, which a smaller batch may already do.

↩ Archetype A (GPT-4/Llama-3), the autoregressive transformer whose serving must stream billions of weights for every generated token, is established in section 1.6.1.1.

Algorithm 10.1 Continuous batching decode scheduler

Require: prefilled waiting queue Q; active set A; KV-cache page allocator; active-token capacity C_{tok}; batched decode kernel
Ensure: generated tokens for completed requests; a continually refreshed active batch

while the service is running **do**
  admit requests from Q while capacity and free KV-cache pages remain; allocate state, append to A
  reserve pages for the next decode step; preempt requests that do not fit
  run one batched decode kernel over the remaining sequences in A
  append each sampled token; extend its KV-cache pages
  remove completed or length-limited sequences from A; return their pages to the allocator
  if HBM pressure exceeds the policy threshold **then**
    move low-priority sequences from A to the paused set; page their KV-cache to CPU DRAM
  end if
  fill freed capacity from Q, or resume paused sequences whose pages now fit
end while

Admitting, completing, and refilling slots on every iteration lets waiting work use capacity released before the longest sequence finishes. The illustrated preemption path adds HBM-to-CPU-DRAM transfers under memory pressure, so recovered capacity must repay both scheduling and state-movement costs. Counted in slot-iterations (figure 10.7), fixed membership leaves only 20 of 32 productive, while continuous admission fills all 32 under the illustrated arrival and capacity assumptions. Prefill and decode cells do not necessarily take equal time. Slot occupancy therefore measures the opportunity; the realized throughput gain depends on how much longer the inserted prefill cells take than the decode steps they share an iteration with.

Figure 10.7: Static vs. Continuous Batching: In static batching (a), completed slots cannot be reused until the longest request finishes, although completed responses may already be released. Continuous batching (b) admits waiting requests into freed slots at iteration boundaries, starting each with a prefill step. Over the eight iterations shown, 20 of 32 slot-iterations do useful work under static batching and all 32 under continuous batching; the hatched idle slots are the capacity recovered when arrivals and cache capacity permit.

10.6.5 Quantifying continuous batching

In a static batch whose output lengths vary widely, a single long response holds every completed slot idle until its final token, wasting a growing share of the batch's sequence-iterations as length variance rises. Continuous batching recovers that stranded capacity by refilling vacated slots at iteration boundaries (figure 10.7). Counting useful and unused sequence-iterations in a fixed batch bounds the recoverable opportunity; arrivals, KV-cache admission, prefill scheduling, and scheduler overhead determine how much of that ceiling a live replica actually captures.

10.6.5.1 The waste function for traditional batching

The 11.25 percent example of one early-finishing sequence generalizes to any mix of output lengths. For this waste model, traditional batching (also called static batching) retains a fixed batch width through all decode iterations until the longest sequence completes. How a batch forms (by count or by time window) and how its membership changes (fixed, or updated at iteration boundaries) are separate choices, so this waste model applies to any fixed-membership batch, including one formed by a dynamic batching window. Assume equal compute cost per sequence-iteration, including padded work for completed requests. For batch size B with output lengths $\{S_1, S_2, ..., S_B\}$, the total compute performed under those assumptions is:

$$O_{\text{traditional}} = B \times S_{\max} \times c_{\text{decode}}$$

where $S_{\max} = \max_i(S_i)$ and c_{decode} is the compute cost per decode iteration per sequence. However, the useful compute is only:

$$O_{\text{useful}} = \sum_{i=1}^{B} S_i \times c_{\text{decode}}$$

Equation 10.4 defines the waste ratio that quantifies the inefficiency:

$$W = 1 - \frac{O_{\text{useful}}}{O_{\text{traditional}}} = 1 - \frac{\sum_{i=1}^{B} S_i}{B \times S_{\max}} = 1 - \frac{\bar{S}}{S_{\max}} \tag{10.4}$$

where $\bar{S}$ is the mean output sequence length. Waste therefore depends entirely on the ratio of mean to maximum output length within the batch. For uniform output lengths ($\bar{S} = S_{\max}$), waste is zero. For highly variable lengths, waste can exceed 50 percent.

10.6.5.2 Serving variable-length outputs in LLMs

Consider a GPT-class model serving four concurrent requests whose prompt and output lengths vary fourfold across the batch. Table 10.6 lists the token counts for each request:

Table 10.6: Concurrent Request Generation Lengths: Four concurrent requests on a GPT-class model with varying prompt and output lengths. The 4× spread in output length (50 to 200 tokens) drives batching waste, because a static batch of four pads every request to the longest, wasting compute on the three shorter requests until the longest finishes generating. Prompt length sets each request's prefill work, and the total is the number of tokens its KV cache holds when it finishes; both matter for admission, while the decode-slot accounting below uses only the output lengths.

Request	Prompt length	Output length	Total tokens
R1	100	50	150
R2	80	200	280
R3	120	100	220
R4	90	150	240

Napkin Math 10.3: Continuous batching waste calculation

Problem: Four LLM requests with unequal output lengths share a four-slot batch. Assume fixed decode-iteration time and no queueing. How much of the batch's slot time does fixed membership leave unused, how much of it can continuous admission hand to new requests, and what happens to the original requests' latencies?

Variables:

- Decode time per iteration (batch of 4): 20 ms
- Maximum output length in batch: 200 tokens (R2)
- Mean output length: $(50 + 200 + 100 + 150)/4 = 125$ tokens

Slot accounting under fixed membership: The batch keeps all four slots until R2, the longest request, finishes.

- Total decode iterations: 200
- Total batch time: 200×20 ms $= 4{,}000$ ms

The waste ratio in equation 10.4, one minus the ratio of mean to maximum output length, measures the slot time that finished requests leave idle:
$W = 1 - \frac{125}{200} = 1 - 0.625 = 37.5\%$
The fixed-width schedule contains $4 \times 200 = 800$ "sequence-iterations" but only 500 are useful. These counts measure executed work only if the kernel continues processing completed slots; otherwise they measure useful work relative to the schedule's available slots.

Continuous admission: Sequences depart at the iteration where they complete, and eligible new requests take their slots without delaying the original requests under the fixed-iteration-time assumption.

- Iteration 50: R1 completes $\rightarrow$ slot freed, new request R5 can join
- Iteration 100: R3 completes $\rightarrow$ slot freed, new request R6 can join
- Iteration 150: R4 completes $\rightarrow$ slot freed, new request R7 can join
- Iteration 200: R2 completes

Response latency, a separate effect: The fixed-membership schedule in this comparison deliberately withholds all four responses until R2 completes its 200 tokens. Static batch membership alone does not require that barrier.

- Completion times with the barrier:
  - R1 completes useful work at iteration 50, but waits until iteration 200 $\rightarrow$ latency = 4,000 ms
  - R2 completes at iteration 200 $\rightarrow$ latency = 4,000 ms
  - R3 completes useful work at iteration 100, but waits until iteration 200 $\rightarrow$ latency = 4,000 ms
  - R4 completes useful work at iteration 150, but waits until iteration 200 $\rightarrow$ latency = 4,000 ms

Latencies with immediate release under continuous admission (assuming no queuing delay):

- R1: 50×20 ms $= 1{,}000$ ms (4× improvement over traditional)
- R3: 100×20 ms $= 2{,}000$ (2× improvement)
- R4: 150×20 ms $= 3{,}000$ (1.33× improvement)
- R2: 200×20 ms $= 4{,}000$ (no improvement for longest request)

Average latency comparison

- Traditional: 4,000 ms (all requests)

- Continuous: (1,000 ms + 2,000 + 3,000 + 4,000) / 4 = 2,500

Result: Under fixed membership, only 500 of the 800 sequence-iterations do useful work. Continuous admission hands the rest to R5, R6, and R7, which start generating at iterations 50, 100, and 150 instead of waiting for a new batch, so the same 200 iterations now advance seven requests rather than four. Releasing each response on completion separately reduces the original requests' average latency by 37.5 percent, numerically matching the unused-slot ratio, and static batching that released responses on completion would obtain the same original-request latencies.

Systems insight: Response release and slot reuse are separate decisions. Removing a response barrier helps the requests already in the batch; only continuous admission recovers the stranded iteration capacity shown in figure 10.7, and that recovered capacity is what raises a replica's throughput.

10.6.5.3 Efficiency regimes for continuous batching

The continuous-batching analysis relates recoverable slot capacity to the gap between a batch's longest and average output lengths. Let $\text{CV} = \sigma/\mu$ be the coefficient of variation of output lengths. For the idealized fixed-width schedule, suppose the longest request in a batch is roughly k standard deviations above the mean. The factor k is an assumption about the batch rather than a property of the workload, because a larger batch draws more samples and pushes its maximum further into the tail. With immediate refill and unchanged iteration time, the corresponding capacity ratio is:

$$\text{Improvement} \approx \frac{S_{\max}}{\bar{S}} = \frac{\mu + k\sigma}{\mu} = 1 + k \cdot \text{CV}$$

where k specifies how far the batch maximum lies above the mean, so the expression is a scenario ceiling on the opportunity from filling otherwise unused slots; scheduler overhead, KV-cache pressure, and refill gaps reduce the recovered fraction.

Table 10.7 applies this ceiling across workload types, including retrieval-augmented generation (RAG) (Lewis et al. 2020). Its CV and k columns are inputs. Its speedup column is an assumed realized fraction held below the $1 + k \cdot \text{CV}$ ceiling. Its waste column is the effective waste that each assumed speedup implies through speedup $\approx 1/(1 - W_{\text{effective}})$. That waste is smaller than the slot waste of equation 10.4, because each speedup sits below its ceiling. Neither the speedup nor the waste column is a measurement. Together they show how much of the ceiling a scheduler must recover to reach a given gain, which is the quantity a load test should measure.

Table 10.7: Continuous Batching Benefit by Workload Type: Higher output-length variance (CV) raises the ceiling on improvement from continuous batching. The k values are assumed, so each $1 + k \cdot \text{CV}$ ceiling is a scenario value, and the speedups are assumed values held below it rather than measurements; the waste column is the effective waste those speedups imply. Workloads with predictable output lengths (code completion) have little to recover, while highly variable workloads (RAG with documents of varying length) have the most.

Workload type	CV	k	Waste (trad.)	Speedup (continuous)
Code completion	0.3	2.5	16.7%	1.2×
Chat (short responses)	0.6	2	33.3%	1.5×
General text generation	1	2.5	50%	2×
Creative writing	1.5	3	64.3%	2.8×
RAG with variable docs	2	2.5	71.4%	3.5×

Continuous batching is most valuable when output lengths are unpredictable, mixed workload types share the same cluster, request volume is high enough to refill vacated slots, and latency SLOs make early completion valuable. It provides minimal benefit when output lengths are uniform, as in classification or embedding generation, when batch sizes are too small for slot reuse to matter, or when request volume is too low to refill vacated slots.

Because arrivals, KV-cache admission, prefill scheduling, and scheduler overhead decide how much vacated capacity is actually recovered, serving teams must compare policies under the same arrival trace, latency objective, and memory budget. The published vLLM system, which pairs iteration-level admission with paged KV allocation, reports 2–4× higher throughput at comparable latency than FasterTransformer and Orca on its evaluated workloads (Kwon et al. 2023). That figure compares whole systems rather than isolating slot reuse, because paged allocation also permits larger active batches.

10.6.5.4 Implementation complexity trade-offs

The same variability that makes continuous batching valuable also makes it harder to operate. Memory management must support overlapping request lifetimes. Fixed-membership batching can reserve cache regions when a batch forms and release them together, whereas continuous admission must track which request owns each region and when it becomes reusable, either through fixed-capacity slots that retain unused reservations or through paging with its block tables. Kernels must also handle variable-length sequences and changing batch membership, often by packing short sequences under shared attention masks.

Scheduler work moves into the decode loop as well. A fixed-membership policy checks once per batch whether it is full or its timeout has expired, while continuous admission tracks completions, free cache capacity, and eligible waiting requests at every iteration boundary. That work must stay small relative to a decode step or it erodes the throughput that slot reuse recovers, so engines keep the per-iteration path short and prepare the next iteration's batch on the host while the accelerator executes the current one. Continuous batching therefore buys its throughput gain with implementation and operating cost, weighed dimension by dimension in table 10.8.

Table 10.8: Traditional vs. Continuous Batching Trade-Offs: Continuous batching provides significant throughput gains for variable-length workloads at the cost of implementation complexity. For uniform-length workloads, the complexity overhead may not justify adoption.

Dimension	Traditional batching	Continuous batching
Implementation effort	Low (standard frameworks)	High (custom scheduler, kernels)
Memory overhead	Fixed allocation	Dynamic + fragmentation mgmt
Scheduler work	Once per batch	Every decode iteration
Debugging complexity	Deterministic behavior	State-dependent, harder to trace
Throughput (variable)	Baseline	Up to $1 + k \cdot \text{CV}$
Throughput (uniform)	Baseline	~1.0× (no improvement)

Because continuous batching carries this implementation cost, some services still run request-level batching, and its barrier can surface as a tail-latency regression whose cause is not obvious. A constructed p99 regression shows how to investigate infrastructure, execution, and serving layers to find that cause. Its latency observations and mitigation outcome are scenario assumptions; only the transfer-time estimates are computed from payload and bandwidth.

Example 10.2: Debugging high p99 latency

Scenario: An LLM serving system shows high tail latency: p50 = 100 ms and p95 = 180 ms maintain SLO, but p99 spikes to $T_{\text{lat,p99}} = 500$ ms against a 200 ms target despite a monitored GPU utilization of 85 percent.

Diagnosis: The system runs across 4 A100 GPUs linked via PCIe Gen4 (32 GB/s per GPU) rather than NVLink. For a 100 MB payload, the bandwidth-only transfer estimate is $T_{\text{comm}}(N) \approx 3.1$

ms over PCIe versus 0.17 ms over NVLink. Collective startup, topology, and contention can add time. The service still runs an older request-level batching server rather than an iteration-level engine, with max batch size 32 and timeout 50 ms, so a batch holds its members until its longest generation finishes. The scenario assumes that 5 percent of batches experience head-of-line queuing delay $L_{\text{lat,queue}}$ as short requests wait behind long generations.

Systems lesson: The latency distribution motivates tracing both candidates, because a transfer paid by every request also amplifies tail queueing near saturation. In this scenario, traces showing stable transfer times too small to explain the additional delay, together with short requests waiting behind long generations, identify the fixed-membership barrier as the immediate cause. Continuous admission removes that barrier; prefill scheduling, cache capacity, and other waits still require measurement. Until the server can be replaced, the following policy reduces output-length variation within request-level batches:

1. **Request classification**: Classify requests by expected output length (short: under 50 tokens, medium: 50 to 200 tokens, long: over 200 tokens)
2. **Differentiated batching**: Limit short batches to 8, medium to 16, long to 32

Grouping requests by expected length shrinks the spread of output lengths inside each batch, which is the quantity equation 10.4 says drives both waste and waiting. The scenario assumes a post-mitigation p99 latency of $T_{\text{lat,p99}} = 185$ ms; a real service confirms it by replaying the same arrival workload and measuring the tail.

10.6.6 Feature-parallel batching for recommendation systems

A single recommendation request may issue thousands of sparse lookups across embedding tables holding billions of entries before the dense ranking head can score a candidate. Grouping requests solely at the dense MLP therefore leaves the service bottlenecked on scattered memory reads and cross-server transfers. While fixed-shape vision inputs allow simple request grouping and autoregressive generation pairs iteration-level admission with KV-cache management (and with prefill-decode scheduling in section 10.8), recommendation systems must also batch around the placement of sparse state. *Sparse embedding lookups* can dominate memory access and communication in an embedding-heavy workload even when *dense matrix multiplications* account for much of the arithmetic. Feature-parallel batching groups those accesses by feature type and destination, which improves lookup locality at each shard and complements request batching in the dense stages.

Recommendation batching starts by grouping work around embedding locality rather than request shape. An incoming request splits across user, item, and context feature types to query sharded embedding tables concurrently, after which a gather barrier concatenates the retrieved representations for the dense MLP ranking head (figure 10.8).

The sparse lookup path and the dense feature path can execute concurrently before feature interaction combines their representations and a ranking head produces final scores. The interaction stage must wait for its required inputs, which puts the slowest embedding shard on the request's critical path.

When sparse embedding access dominates latency, feature-parallel batching groups lookups by feature type and destination across requests. Each embedding server then handles the lookups for its shard across the whole batch. This turns feature sharding into a serving strategy rather than a storage detail, while preserving the association needed to reconstruct each request for ranking:

```
Request 1: [user_id_1, item_ids_1, context_1]
Request 2: [user_id_2, item_ids_2, context_2]
Request 3: [user_id_3, item_ids_3, context_3]

Feature-parallel view:
User embeddings:  [lookup(user_1), lookup(user_2), lookup(user_3)]  → parallel
Item embeddings:  [lookup(items_1), lookup(items_2), lookup(items_3)]  → parallel
Context features: [process(ctx_1), process(ctx_2), process(ctx_3)]  → parallel
```

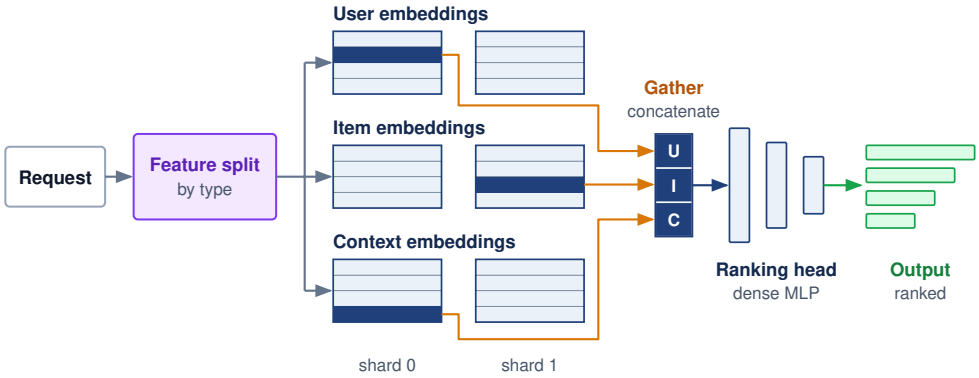

Figure 10.8: Feature-Parallel Batching Pipeline: Requests are batched by feature type (Users, Items, Context) and dispatched to specialized embedding servers in parallel. The retrieved embeddings are then concatenated and processed by the dense ranking head. This architecture enables scaling to trillions of parameters by decoupling embedding storage from dense compute.

```
Then: Combine features per request for ranking
```

Example 10.3: RecSys batching at hyperscale

Scenario: A hyperscale recommendation service handles 10 million QPS, querying 5,000 embedding lookups per request across 1,000 shards.

Diagnosis: The assumed distribution of requests requires 50 million lookups/sec per shard. That rate is the capacity target each shard's lookup service must sustain. Accumulating requests over a 1 ms window forms batches of 50K lookups per shard, creating an opportunity to deduplicate, reorder, or coalesce lookups. Batching alone does not make random addresses sequential.

Systems lesson: Shard-level batching exploits sparse feature locality, but its benefit depends on cache hits, access distribution, and the amount of work each shard receives. The formation window spends part of the 5.2 ms end-to-end latency budget; measured lookup and communication times must fit in what remains alongside the dense stages.

10.6.7 Streaming inference for real-time applications

Holding a live audio chunk or robot control frame in a batching queue until a large batch fills stalls transcription and breaks closed-loop control. Speech recognition, video analysis, and robotics process an ongoing sequence of observations under tight per-frame deadlines, so batching across ready streams helps only while the added waiting time fits each frame's slack.

Streaming inference processes inputs incrementally rather than waiting for the complete recording or interaction. Illustrative input cadences include 10–20 ms audio chunks for speech recognition, 30–60 FPS video, and 100–1000 Hz sensor or control updates in robotics. These rates specify how often work arrives, not a universal inference deadline. Buffering, frame selection, and batching must follow the application's latency and freshness requirements.

A streaming service must sustain the arrival rate while meeting the latency budget for each accepted input. The end-to-end processing path includes:

$$T_{\text{streaming}} = T_{\text{capture}} + T_{\text{transfer}} + L_{\text{lat,infer}} + T_{\text{action}}$$

where the sum must fit the end-to-end latency budget. A pipelined implementation must also sustain the input arrival rate; its slowest stage must process frames at least as quickly as they arrive. The end-to-end budget may span several frame intervals. A speech pipeline with 20 ms frames and

a 100 ms budget can hold at most five frames of delay, so batching ready frames from independent streams helps only while the added waiting and service time fit that slack.

10.6.8 Adaptive batching strategies

A batching window and maximum batch size tuned for daytime peak traffic either hold requests waiting too long when traffic drops at night or cap throughput too early when a flash crowd arrives. Because fixed batching parameters break as arrival rates and request mixes shift, production systems adapt their batching windows and size limits to live conditions.

Traffic-adaptive batching adjusts the batch window based on arrival rate:

$$T_{\text{window}} = \min\left(T_{\max}, \frac{B_{\text{target}}}{\lambda_{\text{arr}}}\right)$$

When traffic is high, the window shrinks because the target batch size fills quickly. When traffic is low, the window extends but is capped to bound batch-formation delay. This cap does not bound the separate wait for a busy replica.

SLO-adaptive batching takes a complementary approach, monitoring latency percentiles and adjusting parameters aggressively:

```
if P99_latency > 0.9 * SLO:
    reduce B_max by 20%
    reduce T_window by 20%
elif P99_latency < 0.5 * SLO:
    increase B_max by 10%
    increase T_window by 10%
```

The feedback loop is a tuning heuristic, not a guarantee of latency headroom. Reducing batch size can shorten execution while also reducing service capacity and increasing queueing under heavy load. Feedback delay and noisy percentile estimates can cause oscillation, so the controller must be evaluated with the intended arrival trace and overload policy.

Request-aware batching adds another dimension by considering request characteristics. For LLMs, prompt length can guide padding and prefill budgets, while estimated output length can guide grouping. Prediction errors and priority choices must be included when evaluating tail latency.

Production runtimes expose these controls directly. NVIDIA Triton Inference Server, for example, caps batches with `max_batch_size`, bounds formation waiting with `max_queue_delay_microseconds`, and accepts `preferred_batch_size` values for particularly efficient execution profiles; priority queues and overload policy remain separate controls. Under a load test, larger batches absorb rising traffic until offered load approaches execution capacity, after which completions fall below arrivals and the growing queue needs an admission or scaling response rather than a larger batch.

The batching strategies to this point take each request's work as given, apart from a cap on generation length. Reasoning-heavy inference makes that budget an explicit quality-cost decision. The service may allocate additional tokens or candidate evaluations when the expected benefit justifies their latency and capacity costs.

10.6.9 Test-time compute and the per-request budget

Letting a model emit hundreds of serial chain-of-thought (CoT) tokens or sample dozens of candidate solutions per query can turn a 100 ms interactive reply into a multi-second job that pins KV-cache memory and drains replica capacity. Test-time scaling buys answer quality by spending more compute on each query at inference time, which moves the scaling curve out of the training cluster and into the request path. Every increment of quality bought this way is paid on every request for the lifetime of the deployment, the recurring cost that the serving cost dominance principle (17) already prices. Whether extra generated tokens, internal CoT steps, or search branches actually improve quality on a given task is an empirical question. Capability studies describe abilities that appear abruptly with scale (Wei et al. 2022), while later analysis finds some apparent jumps to be artifacts of the metric chosen (Schaeffer et al. 2023). A serving budget should therefore rest on measured quality gains for the workload.

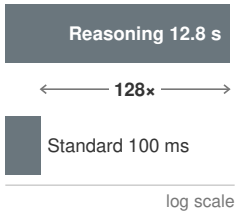

Reasoning through 128 serial tokens turns a 100 ms answer into 12.8 s.

From a serving-systems perspective, the relevant shift is the work issued per request. Additional reasoning consumes accelerator time and can extend both latency and KV-cache residency. Increasing test-time work does not necessarily change the bottleneck of each decode step. It instead requires a **per-request compute budget**, a limit on the generated tokens, candidate solutions, or search steps a request may consume, stated in both latency and capacity. The budget binds when serial reasoning tokens exhaust a request's latency target, or its share of fleet capacity, before additional reasoning stops improving the answer. A harder task may justify more work, but additional work does not guarantee a correct answer. Adaptive batching decides when requests run together; test-time compute allocation decides how much work each request is allowed to issue.

↩ Test-time scaling, the lifecycle regime that spends compute per request rather than per training run, is established alongside the empirical scaling laws in section 1.3.3.1.

Napkin Math 10.4: Scaling reasoning depth

Problem: Calculate the latency impact of a model that uses 128 "Thinking Tokens" to solve a complex math proof vs. a standard answer.

1. **Standard response**: 1 token answer = 100 ms.
2. **Reasoning response**: 128 tokens of internal search/CoT before the answer.
3. **The latency**: 128 $\times$ 100 ms = 12.8 seconds.

Systems insight: The per-token time is a round figure chosen for easy arithmetic, more than ten times the batch-one decode step of the reference 70B replica in section 9.6.3. What the calculation teaches does not depend on it. Serial reasoning tokens multiply latency by the ratio of reasoning tokens to answer tokens, whatever each token costs, so a faster replica shortens both responses in proportion and leaves the 128-fold expansion intact. Against a direct answer of many tokens rather than one, the ratio shrinks accordingly; against a one-token answer it is largest. A search step that evaluates several candidates in parallel adds compute and KV residency without adding the same serial latency, which is why reasoning budgets must be stated in both time and capacity. At fleet scale, compute allocation weighs the measured quality gain against both the request deadline and the capacity the additional work consumes.

Variable work per request closes the batching story: continuous and adaptive batching handle variation the system observes, while test-time compute is variation the system chooses. A serving policy therefore sets each request's budget from task difficulty, latency tolerance, and fleet load rather than applying a fixed decode budget to every prompt. With the batching mechanism and the per-request compute budget both defined, the remaining decision is which combination a given workload should run.

10.6.10 Selecting a batching strategy

Each batching mechanism trades scheduler overhead against slot utilization, so a fixed-shape vision classifier needs a different policy than an autoregressive chat endpoint or a sparse recommendation ranker. Selecting a batching strategy starts by matching the mechanism to the model's work shape, traffic variance, and where the request path spends its latency budget.

Checkpoint 10.3: Batching strategy trade-offs

Verify your understanding of different batching mechanics:

- ☐ **Static vs. continuous**: for a vision service with highly uniform input sizes and predictable traffic, select static batching or continuous batching.
- ☐ **Waste ratio**: explain when refilling completed slots can eliminate the static-batch padding waste ($1 - \bar{S}/S_{\text{max}}$), and why low arrivals or KV-cache pressure prevent perfect slot reuse.
- ☐ **Adaptive batching**: explain when shrinking the formation window reduces waiting without sacrificing throughput, and why it cannot resolve overload when arrivals exceed service capacity.

- ☐ **RecSys bottleneck**: in feature-parallel batching for RecSys, identify whether the primary fleet-stack bottleneck is the GPU's Tensor Cores or the distributed Embedding Store's memory bandwidth.
- ☐ **Batching limit**: explain why free HBM, not Tensor Core compute, often caps the achievable batch size for long-context LLM serving.

Noncompute stages such as host serialization, queuing, and network routing consume a substantial fraction of the latency budget traced in figure 10.1, so batch formation and execution must be budgeted across the full pipeline rather than against accelerator runtime alone. The decision tree's first branch names chunked prefill, which splits a long prompt into chunks interleaved with ongoing decode steps (section 10.8.2). The following questions identify candidate policies; memory admission, arrivals, and the remaining latency budget determine their usable settings:

```
Is it autoregressive text generation?
|- Yes → Continuous batching with chunked prefill
`- No → Is it real-time streaming speech/video/robotics?
        |- Yes → Streaming pipeline with minimal or no batching
        `- No → Do embedding lookups dominate latency?
                |- Yes → Feature-parallel batching (RecSys)
                `- No → Dynamic batching with adaptive parameters
```

Each resulting strategy has a typical batch range, a constraint that binds it, and a parameter it tunes (table 10.9). Vision and long-prompt prefill can emphasize compute utilization, small-batch decode can emphasize weight bandwidth, and sharded recommendation can emphasize embedding access, while speech introduces deadlines tied to the input stream. Each row is a starting point for measurement against both useful throughput and the latency budget, not a fixed bottleneck or batch-size requirement for every model in a category.

Table 10.9: Batching Strategy by Model Type: Each model class pairs a batching policy (from static or dynamic batching in vision to continuous batching in LLM decode and feature-parallel batching in recommendation) with its dominant constraint and the parameter that tunes it. Typical batch sizes are operating ranges set by latency budgets and KV-cache capacity; they may fall below or above the overhead-amortization knees of notebook 10.2, which mark where fixed costs stop dominating rather than where batching stops paying.

Model type	Strategy	Typical batch size	Key constraint	Key parameters	Tuning goal
Vision (CNN)	Static/Dynamic	32–256	GPU compute	Batch size; window and max batch	Balance latency vs. throughput
LLM (prefill)	Dynamic	1–64	GPU compute	Window, max batch	Balance latency vs. throughput
LLM (decode)	Continuous	Tens–hundreds (KV-limited)	Memory bandwidth	Chunk size, max batch	Minimize decode latency variance
RecSys	Feature-parallel	1000–10000s	Embedding lookup	Accumulation window	Match embedding shard capacity
Speech	Streaming	1	Real-time	Pipeline depth	Meet real-time deadline

Batching policy and memory admission are coupled. In a full-attention deployment that retains active KV state on the GPU, every admitted sequence consumes cache capacity as its context grows. A schedule with room for more decode work can therefore run out of memory before it reaches its execution limit. Managing that state is the next step in determining useful concurrency.

10.7 Memory and Decode-Time Management

As a full-attention LLM generates a 2,000-word essay, each new token attends to representations of earlier tokens. The KV cache retains the keys and values needed for that attention, with a per-request footprint that grows as generation proceeds. Large contiguous reservations can leave HBM reserved

but unused or create gaps too small for another reservation. Admission can then fail even though the stored attention state occupies less than the available capacity.

Definition 10.2: KV cache

KV cache is a per-request LLM inference memory buffer that stores key and value attention tensors for processed prompt and generated tokens. It avoids recomputing those tensors for the prefix; each new query still performs attention against the cached state.

1. **Significance**: Compared with recomputing full-prefix attention for each new token, caching changes the attention work from $\mathcal{O}(S^2)$ to $\mathcal{O}(S)$ in sequence length for full attention with fixed model dimensions. Retaining that state costs memory proportional to context length and the number of active sequences, potentially exceeding the weight footprint. Section 10.7.1 sizes it per token and per replica.
2. **Distinction**: Model weights are shared across requests using the same model, while each sequence has its own logical KV state. Physical storage can still be shared for compatible common prefixes. The scheduler must budget for distinct retained state and its growth, rather than assume every request needs a wholly separate copy or that sharing removes the capacity limit.
3. **Common pitfall**: A frequent misconception is that reducing model size proportionally reduces serving memory. In practice, the KV cache dominates HBM at long contexts, so a quantized 70B model can still be capacity-bound at batch sizes the weights alone would easily admit.

Two classes of techniques meet at this boundary. KV-state capacity techniques, such as PagedAttention and prefix caching, decide where attention state lives and how much fragmentation the memory manager tolerates. Decode-time latency techniques, such as speculative decoding, change how much target-model work is needed per emitted token. Speculative decoding is not KV-cache compression; it trades extra draft-model state, target-model verification, and scheduler complexity for lower TPOT under the same memory budget. Capacity accounting and adapter locality then determine how retained KV state and task-specific weights share the remaining memory budget.

Most of the calculations that follow share one reference replica, the 70B model on one eight-GPU H100 node that fleet efficiency measured as its serving replica, with FP16 weights split eight ways; the few that leave it say so. The attention layout matters even more, because it sets the bytes cached per token. Grouped-query attention (GQA) with eight KV heads, the layout of the Llama-family 70B models, caches one eighth the state of a 64-head multi-head attention baseline, because groups of query heads share each cached key and value head (section 10.7.5.2). Each calculation therefore names its hardware and attention layout, and numbers from different calculations compare only after conversion to the same configuration.

↩ The reference serving replica, a 70B model on one eight-GPU H100 node, is established in section 9.1.

10.7.1 The KV cache wall and memory capacity

Even when a 4-bit quantized 70B model occupies only 35 GB of an 80 GB H100, admitting a second 128K-token request pushes the combined weight and KV-cache footprint past 80 GB and exhausts HBM before execution resources saturate. Increasing batch size in LLM serving therefore hits the KV cache wall long before compute binds. Figure 10.9 leaves the reference replica to show the wall on a single H100. It combines an idealized 35 GB weight payload for a quantized 70B model with its full-attention, eight-KV-head GQA state. The budget reads the nameplate label in decimal gigabytes, slightly less than the device's registered byte capacity, which keeps the plotted wall conservative; the node estimator below uses the registered byte capacity. A batch of 8 reaches the memory boundary at roughly 17K tokens per sequence; quantization metadata, runtime buffers, and other allocations make the plotted boundary an optimistic sizing limit.

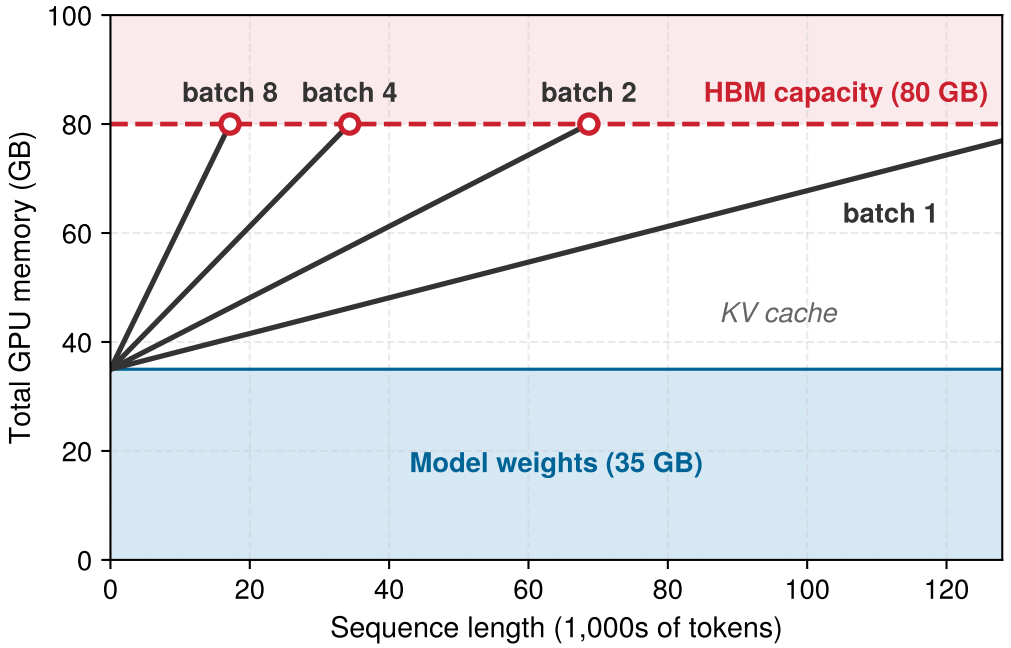

Figure 10.9: The KV Cache Wall: Total GPU memory usage for a 4-bit quantized 70B model (35 GB weights) on an 80 GB H100. While the model fits easily at small context lengths, the KV cache (sloped lines) eventually consumes all remaining HBM. At 128K context, a batch size of 2 is physically impossible on a single GPU (35 GB + 84 GB exceeds 80 GB), forcing the system to either reduce batch size, change the state representation, or distribute the state across devices.

The visualization explains why weight fit alone does not establish serving capacity. A long-context request can consume nearly all remaining memory in the illustrated configuration, leaving little space for other active sequences. Sharding can add headroom, but its communication cost must be weighed against alternatives such as lower concurrency or a smaller cache representation. The effect of a 128K context depends on the model's cache dimensions and precision, not on context length alone.

The capacity model behind the figure starts from the cost of one cached token. Every layer stores a key vector and a value vector for each KV head, so with N_L layers, H_{KV} KV heads, head dimension d_{head}, and s_{elem} bytes per element, equation 10.5 gives the bytes cached per token:

$$m_{\text{token}} = 2 \times N_L \times H_{\text{KV}} \times d_{\text{head}} \times s_{\text{elem}} \tag{10.5}$$

A replica's KV footprint is that cost multiplied by every token it holds, $M_{\text{KV}} = B \times S \times m_{\text{token}}$, so the largest admissible batch at context S is the HBM left after weights and runtime allocations divided by $S \times m_{\text{token}}$. Each serving lever acts on a different factor. Grouped-query attention lowers H_{KV}, KV quantization lowers s_{elem}, a shorter admitted context lowers S, and paging returns HBM lost to fragmentation without changing the per-token cost. The notebook below turns this budget into an explicit batch-size limit for a production node.

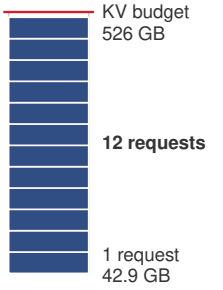

A dozen 128K-token requests exhaust the whole KV budget.

Napkin Math 10.5: KV-cache capacity estimator

Unlike the conservative decimal budgets used in the illustrative comparisons, this estimator retains the registered byte capacity when summing the node's HBM. Converting its display units does not change that byte count.

Problem: A deployment serves an illustrative 70B-class GQA model with 128K context support (FP16 weights ≈ 141.2 GB) on an 8× H100 node (640 GB total HBM). The goal is to determine the maximum batch size for a context length of 128K tokens.

Formula: Bytes per token come from equation 10.5, and

$$M_{\text{total}} = M_{\text{weights}} + M_{\text{KV}}, \qquad M_{\text{KV}} = B \times S \times m_{\text{token}}$$

Parameters:

- $N_L = 80$, $H_{\text{KV}} = 8$ for the assumed GQA architecture, $d_{\text{head}} = 128$.
- $s_{\text{elem}} = 2$ bytes (FP16).
- Context $= 131{,}072$ tokens.

Step 1: Calculate memory per token.

$$m_{\text{token}} = 2 \times 80 \times 8 \times 128 \times 2 = 327{,}680 \text{ bytes} \approx 0.33 \text{ MB/token}$$

Step 2: Calculate cache per request.

$$131{,}072 \text{ tokens} \times 0.33 \text{ MB/token} \approx 42.9 \text{ GB/request}$$

Step 3: Determine max batch size. Available Memory for KV = 687.2 GB (Total) - 141.2 GB (Weights) - 20 GB (System) = 526.0 GB.

$$\text{Max Batch} = \left\lfloor \frac{526.0}{42.9} \right\rfloor = 12$$

Systems insight: Even with GQA, whose shared key-value heads the 8-head count in this calculation already reflects, each 128K-context request consumes tens of gigabytes of KV state. This configuration therefore admits only a limited number of concurrent long-context requests. PagedAttention can recover allocation waste, but cannot shrink the attention tensors counted here. Increasing concurrency beyond that limit requires changing the memory budget, cache representation, or placement of state.

10.7.2 The fragmentation problem

An allocator that reserves a contiguous HBM slab for every sequence's maximum permitted length strands KV-cache memory in unused reservations and noncontiguous gaps (Kwon et al. 2023). Traditional contiguous allocation creates two distinct forms of this waste.

Internal fragmentation wastes memory within each allocation. Preallocated memory management reserves room at admission for the prompt plus the permitted output, so a request allowed 4,096 output tokens holds that room even if generation stops after 50. The unused reservation cannot serve another request until it is released. If the output reservation holds 4,096 tokens but the average output uses 100 tokens, 97.6 percent of that output reservation goes unused. Prompt state consumes additional space; this calculation does not give the unused fraction of the combined prompt-and-output allocation.

External fragmentation compounds this problem across allocations. As sequences complete and new ones start, memory becomes fragmented into noncontiguous free blocks. Even with sufficient total free memory, no single block may be large enough for a new maximum-length allocation.

Consider a simplified example with 8 memory slots and maximum sequence length of 4. Fixed-size reservations first create internal fragmentation:

```
Time 0: Allocate Seq A (slots 0-3), Seq B (slots 4-7)
        [A][A][a][a][B][B][B][b]
        lowercase = reserved but unused
```

External fragmentation can arise when the allocator also permits smaller contiguous reservations. In a later state, A and B have finished and released their slots, and two shorter sequences, C and D, hold two-slot reservations, leaving separated gaps:

```
Time 1: Active Seq C and Seq D leave two 2-slot gaps
```

```
        [ ][ ][C][C][ ][ ][D][D]

Time 2: Try to allocate Seq E (needs 4 contiguous slots)
        [ ][ ][C][C][ ][ ][D][D]  <- Total free slots = 4, largest block = 2

Result: enough total free capacity exists, but no contiguous block is large enough.
```

Under variable sequence lengths, the two effects compound. The vLLM authors measured roughly 60 to 80 percent of reserved KV-cache memory going unused in the prior serving systems they studied, capacity that could otherwise have admitted more sequences (Kwon et al. 2023).

10.7.3 PagedAttention

The fragmentation example exposes the allocator failure: enough KV-cache capacity can exist in aggregate while no contiguous block is available for the next sequence. PagedAttention fixes that serving problem by treating KV-cache storage like virtual memory rather than a single preallocated slab (perspective 10.1 and definition 10.3).

Systems Perspective 10.1: Analogy: The inefficient hotel

Consider a hotel storage room with 100 lockers. Each guest may bring up to 10 suitcases, but currently every guest has only one suitcase.

Under maximum-size reservation, the manager reserves 10 lockers for each guest. Ten guests exhaust the 100-locker budget while occupying only 10 lockers; the other 90 remain reserved but unused.

With on-demand allocation, each suitcase receives any free locker, and a directory records which lockers belong to each guest. At the current one-suitcase occupancy, the room can accommodate 100 guests. Additional suitcases still consume real capacity; allocation does not remove the storage limit.

Definition 10.3: PagedAttention

PagedAttention is an LLM serving memory management technique that applies virtual memory principles to KV cache allocation, storing attention states in noncontiguous, fixed-size physical blocks.

1. **Significance**: It eliminates external fragmentation within the fixed-block pool and bounds internal waste to the unused portion of each sequence's final block. Recovering capacity from contiguous preallocation permits larger concurrent batches; the throughput gain depends on sequence lengths, demand, and the serving bottleneck.
2. **Distinction**: A contiguous allocation must find a sufficiently large free region; supporting later growth may require reserving extra space or relocating the allocation. PagedAttention uses a block table to map logical sequence indices to physical memory pages, allocating blocks as the sequence grows without requiring adjacent free space.
3. **Common pitfall**: A frequent misconception is that PagedAttention speeds up *individual token math*. In reality, it is a capacity optimization: it improves *system-level efficiency* by allowing larger batch sizes, though it adds a small indirection overhead for the block-table lookups.

In vLLM (Kwon et al. 2023), this paged memory layout divides the KV cache into fixed-size pages, using 16-token blocks in the original paper's default configuration, and allocates them on demand. Whereas contiguous reservation strands non-adjacent free slots when earlier sequences finish, each sequence's block table maps logical token pages to scattered physical HBM slots as generation proceeds (figure 10.10). The block size fixes how many tokens each page holds, so a sequence's pages need not form one contiguous allocation run.

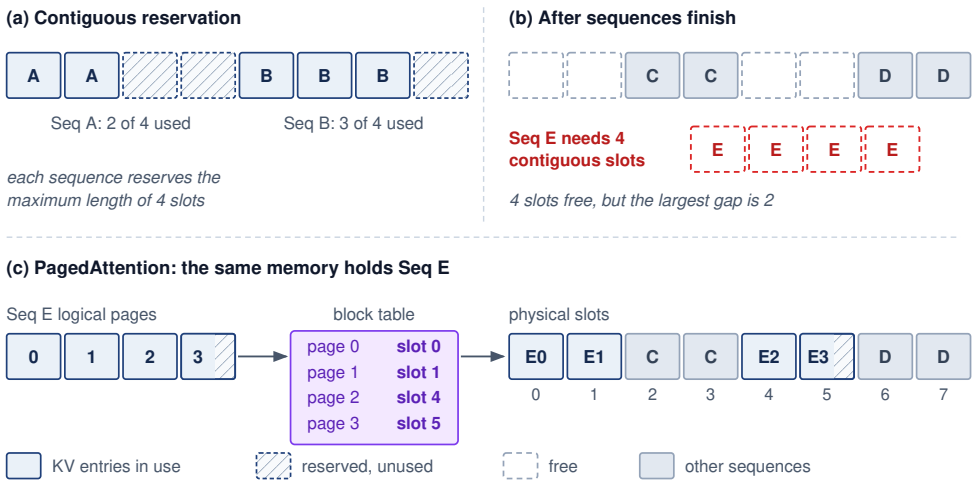

Figure 10.10: KV Cache Fragmentation and the PagedAttention Solution: The panels replay the eight-slot example from the text. (a) Contiguous reservation holds the maximum length of four slots per sequence, so the slots Seq A and Seq B never fill sit idle (internal fragmentation). (b) After sequences finish, four slots are free but the largest contiguous gap is two, so Seq E, which needs four contiguous slots, cannot be admitted (external fragmentation). (c) PagedAttention maps Seq E's logical pages through a block table onto any free physical blocks, so the same memory now holds it; only the unfilled tail of its final block remains unused.

Beyond the fragmentation bounds in the definition, paging gives two further capacity advantages. Sequences grow page by page without worst-case preallocation, and common prefixes across requests can share physical pages directly in memory. Together, these properties turn KV cache capacity into a schedulable resource rather than a fixed per-request reservation. The price is indirection, because custom attention kernels must gather keys and values from noncontiguous blocks through each sequence's block table. Recovered capacity also raises throughput only when demand fills it. Allocator occupancy is not a throughput measurement, since throughput still depends on decode work, prefill interference, and the latency budget.

10.7.4 Prefix caching

Recomputing the same 2,000-token system prompt, few-shot example block, or shared document on every arriving request burns prefill compute and fills HBM with identical copies of the same key-value tensors. Because many LLM workloads prepend those shared prefixes across thousands of concurrent queries, prefix caching stores the prefix's KV-cache blocks once and shares them across compatible requests. Reuse requires the same token prefix and the same model computation, including any task-specific weight adapter and other inputs that affect the cached state. The cache key must distinguish those conditions, and the service must restrict sharing to the permitted users or tenants; matching visible text alone is insufficient. In valid block-level sharing (figure 10.11), logical tables for independent requests map common prefix tokens to physical blocks B0–B5, skipping their repeated prefill work while storing request-specific suffix state separately.

With PagedAttention, prefix caching uses copy-on-write semantics, in which shared blocks stay shared until a request must write into one, and that request then receives a private copy. Three requests that share a system prompt map their first six logical pages to the same physical blocks:

```
System prompt → Physical blocks [0, 1, 2, 3, 4, 5]

Request A page table: [0, 1, 2, 3, 4, 5, 10, 11]  <- shares prefix blocks
Request B page table: [0, 1, 2, 3, 4, 5, 12, 13, 14]  <- shares prefix blocks
Request C page table: [0, 1, 2, 3, 4, 5, 15]  <- shares prefix blocks
```

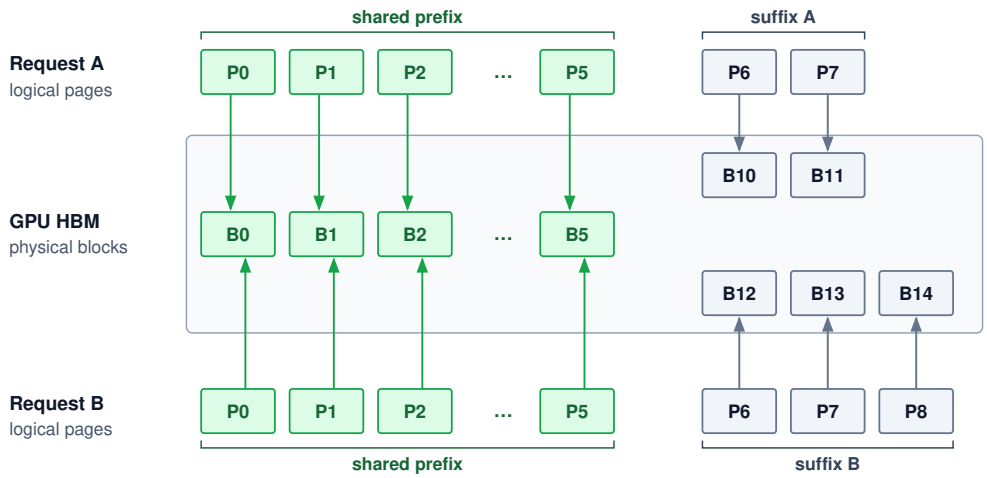

Figure 10.11: Prefix Caching via Block Sharing: PagedAttention enables efficient prefix caching by allowing multiple sequences' block tables to point to the same physical blocks for shared content. In this example, the system prompt is stored in blocks 0–5. Request A and Request B map their first six logical pages to these same physical blocks, storing only their unique suffixes in new blocks. Sharing reduces KV memory and repeated prefill computation for workloads with a common context.

All three requests reference the same physical blocks for the system prompt, and their distinct suffixes occupy new blocks. If a shared prefix ends partway through a block, appending a different suffix writes into a shared block, which is the copy-on-write case. In a chatbot deployment where thousands of concurrent users share a long system prompt, sharing those prefix blocks cuts total KV-cache memory by roughly four-fifths (notebook 10.6).

Napkin Math 10.6: Prefix caching at scale

Scenario: A chatbot service uses a 2,000-token system prompt and serves 1,000 concurrent users. The illustrative model uses multi-head attention with 80 layers, 64 KV heads of dimension 128, and FP16 state. These full-width KV assumptions differ from the grouped-query attention of Llama-70B.

Without prefix caching:

- KV cache per user: $2000 + 500$ (avg response) $= 2{,}500$ tokens
- Total KV cache: 2,500 tokens $\times 1000 \times 2 \times 80 \times 8192 \times 2 = 6.6$ TB

With prefix caching:

- Shared prefix: 2,000 tokens (once)
- Unique per user: 500 tokens
- Total: $(2000 \times 1) + (500 \times 1000) = 502{,}000$
- Memory: $502{,}000 \times 2 \times 80 \times 8192 \times 2 = 1.3$ TB

Result: Prefix caching yields a 79.9 percent memory reduction in KV cache, a roughly 5× capacity ratio for this fixed request population. Actual admitted concurrency also depends on other memory use, generation lengths, and latency limits.

Systems insight: Prefix caching pays off when the prefix hit rate is high because many requests share a long prefix. Table 10.10 uses illustrative hit-rate envelopes to show where the technique changes serving capacity and where it does not; production sizing must derive the distribution from a request trace.

Table 10.10: Illustrative Prefix-Caching Sensitivity: Assumed prefix hit rates and resulting KV-cache memory savings for three workload shapes. The values demonstrate sensitivity rather than population-wide norms; a serving team must measure exact-prefix reuse and prefix length on its own traffic.

Workload	Prefix hit rate	Memory savings
Chatbot (same system prompt)	95%+	70–80%
Document QA (same doc)	80–90%	50–70%
General API (diverse)	20–40%	10–30%

10.7.5 KV cache compression and attention layout

Even after paged allocation (section 10.7.3) eliminates unused reservations and prefix sharing (section 10.7.4) deduplicates shared prompts, unique active tokens still consume enough HBM per token to cap concurrency at small batch sizes. Shrinking that per-token footprint requires reducing either the bytes stored per element or the number of key-value heads cached per layer. The calculation here starts from a 64-KV-head multi-head attention (MHA) baseline so the grouped query attention subsection can isolate the head-count reduction as a separate architectural lever. For a 70B parameter model with 80 layers, 64 heads, head dimension 128, and sequence length 4096 in FP16, the KV cache requires:

$$\text{KV cache} = 2 \times 80 \times 64 \times 128 \times 4096 \times 2 \approx 10.7 \text{ GB (FP16)}$$

A single request's KV cache alone consumes a substantial fraction of the H100's 80 GB of HBM. For a batch of 8 concurrent requests, the KV cache would require 85.9 GB, leaving no room for model weights or runtime state. Quantization reduces bytes per cached element, while architectural optimization can reduce the number of cached elements. These complementary levers change different factors in the capacity calculation.

10.7.5.1 KV cache quantization

Weight-only quantization leaves the KV cache untouched. KV cache quantization instead stores cached keys and values in INT8, FP8, or even INT4. Storing the cache payload in INT4 reduces it to approximately 2.7 GB per request, a 4× reduction before quantization metadata and any higher-precision residual state. Recovered memory can admit more sequences, but throughput depends on kernel support, conversion costs, and the latency target. Quality must also be evaluated for the intended contexts.

Cached keys and values also have different distributions from the model weights and from each other. Keys tend to carry outlier magnitudes, entries far larger than the rest, that persist in the same few hidden channels across tokens. Quantizing keys per token would force every channel in a token's key vector to share one wide scale determined by those hot channels, compressing the remaining channels into only a few quantization levels. Grouping key elements along the token axis within each channel isolates each outlier channel under its own scale. Values lack that fixed channel pattern and can be quantized per token across channels, which lets each new token's value vector be quantized immediately upon generation without rescaling earlier tokens.

10.7.5.2 Grouped-query attention

Grouped query attention (Ainslie et al. 2023) changes the number of key-value heads cached per layer, while quantization changes bytes per cached element. GQA is a model architecture choice established through training or checkpoint adaptation, not a serving switch that discards heads from an unchanged model. Holding the other cache dimensions fixed, the representative comparison of 64 KV heads under MHA with 8 under GQA gives an 8× reduction:

$$\text{KV cache (GQA)} = 2 \times 80 \times 8 \times 128 \times 4096 \times 2 \approx 1.3 \text{ GB (FP16)}$$

At 1.3 GB, the GQA cache occupies less memory than the 10.7 GB required by MHA in this comparison. Combining GQA with INT8 KV cache quantization reduces the payload further,

but per-request size still grows with context length. The architectural and numerical changes each require quality evaluation; recovered capacity permits additional sequences only while their execution remains within the service budget.

Napkin Math 10.7: The precision dividend

The precision dividend is the gain in maximum batch when lower precision frees memory for KV cache. Computing it first with the 64-head MHA baseline, whose requests carry eight times more KV state than the reference replica's eight-head GQA cache, and then with the replica's own GQA cache separates what the attention layout changes from what precision changes.

This comparison uses a conservative decimal HBM budget of 80 GB per GPU, matching the nameplate label rather than the registered byte capacity used by the node estimator.

Problem: A team serves a 70B model on 8× H100 GPUs. Weights in FP16 consume 140 GB (17.5 GB/GPU). KV cache at FP16 consumes 10.7 GB per request. How does quantizing weights to INT4 and KV cache to INT8 change the maximum batch size?

Before optimization (all FP16):

- Weights: 17.5 GB/GPU
- Available for KV cache: 80 GB − 17.5 GB = 62.5 GB/GPU
- KV cache per request: 10.7 GB ÷ 8 ≈ 1.34 GB/GPU
- Maximum batch size: ⌊ 62.5 GB divided by 1.34 GB ⌋ ≈ 46 requests

After optimization (INT4 weights, INT8 KV cache):

- Weights: 17.5 GB $\times(4/16) = 4.38$ GB/GPU (INT4)
- Available for KV cache: 80 GB − 4.38 GB = 75.62 GB/GPU
- KV cache per request (INT8): 5.4 GB ÷ 8 ≈ 0.67 GB/GPU
- Maximum batch size: ⌊ 75.62 GB divided by 0.67 GB ⌋ ≈ 112 requests

The reference replica's GQA cache (same budgets, eight KV heads):

- KV cache per request at FP16: 1.34 GB, or 0.17 GB/GPU, giving a maximum batch of about 372 requests with FP16 weights
- KV cache per request at INT8: 0.08 GB/GPU, giving about 901 requests with INT4 weights, a dividend of 2.4×

Systems insight: The two layouts yield the same multiplicative dividend, because each maximum batch is free memory divided by per-request cache, and precision changes both quantities by the same factors whatever the head count. The layout sets the absolute count, and at either precision the MHA replica admits roughly an eighth as many requests. In both layouts, compressing the cache does more for admission than compressing the weights, since halving per-request state doubles the count while INT4 weights free only the weight share of each GPU. At serving scale the recovered headroom is an admission budget, not yet throughput. The scheduler converts it into completed requests only while the larger batch still meets TPOT, and the counts here omit runtime allocations and headroom, so they are optimistic upper bounds on what admission control should allow.

10.7.6 Speculative decoding

When a lightly loaded replica serves one or a few interactive streams, batching has no other requests to share each weight read with, so every generated token pays the full latency of streaming the target model's weights from HBM. While KV compression and GQA shrink the state each decode step reads, speculative decoding raises the useful work each step performs within a single request. A small draft model proposes K tokens, and the target model verifies all of them in one parallel forward pass (Leviathan et al. 2023; Chen et al. 2023). That pass streams the target's weights once, as a single decode step does, so it processes K candidate positions for roughly the weight traffic of one

token. This is the reuse that batching buys, drawn from one request's future tokens instead of from other requests, and it moves the decode operating point up the bandwidth slope toward the ridge.

On a shared time axis (figure 10.12), standard decoding uses four target passes to emit four tokens. In the speculative round, the draft model proposes T_1 through T_4 sequentially and the target evaluates those positions in one pass; three drafts are accepted, and the first rejected token, T_4, is replaced by T_4' sampled from the target. The shorter timeline holds only when drafting, verification, and correction together cost less than the target passes they replace. If too many drafts are rejected, the target streams its weights and keeps little of the work; if the draft model is too expensive, its own sequential passes consume the time that one-pass verification saved.

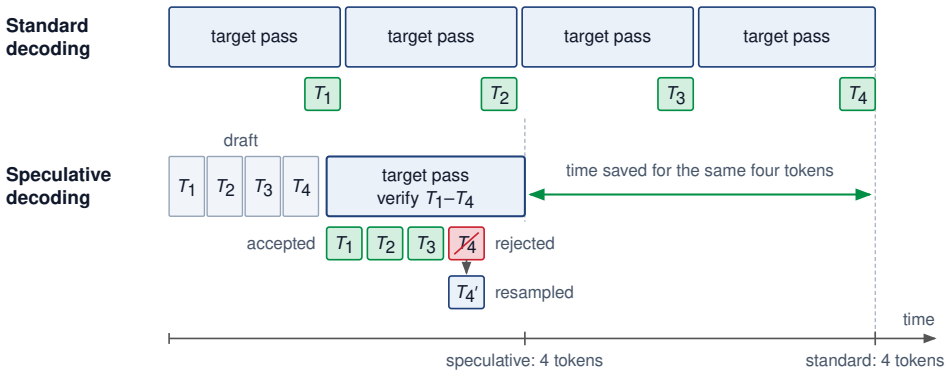

Figure 10.12: Speculative Decoding Process: Instead of generating tokens sequentially with the large target model, a small draft model proposes a sequence of K tokens. The target model verifies all K tokens in a single parallel forward pass. Drafted tokens that pass the acceptance test are emitted together; the first rejected token triggers a corrected sample and discards the remaining draft.

Algorithm 10.2 makes the acceptance test precise. After the single target pass, each drafted token is accepted with probability $\min(1, p_\theta/q_\phi)$, and a rejection resamples from the normalized positive residual $(p_\theta - q_\phi)_+$, so the emitted sequence is distributed exactly as the target model would have produced it unaided. That exactness is what lets a serving system trade draft-model compute for latency without changing the output distribution. Approximate verification rules, which accept drafts the target merely ranks highly, do not inherit the guarantee.

Algorithm 10.2 Speculative decoding

Require: target distribution $p_\theta(\cdot \mid \cdot)$; draft distribution $q_\phi(\cdot \mid \cdot)$; prefix $y_{1:t}$; draft length K
Ensure: one or more new tokens distributed as p_θ

for $i = 1$ to K **do**
  draft $\hat{y}_{t+i}$ from $q_\phi(\cdot \mid y_{1:t}, \hat{y}_{t+1:t+i-1})$
end for
run p_θ once on the prefix plus all K drafts to obtain target distributions for the drafted positions and final next-token position
for $i = 1$ to K **do**
  $a_i \leftarrow \min\left(1, \frac{p_\theta(\hat{y}_{t+i} \mid y_{1:t}, \hat{y}_{t+1:t+i-1})}{q_\phi(\hat{y}_{t+i} \mid y_{1:t}, \hat{y}_{t+1:t+i-1})}\right)$
  accept $\hat{y}_{t+i}$ with probability a_i
  if $\hat{y}_{t+i}$ is rejected **then**
    sample one corrected token from the normalized positive residual $(p_\theta - q_\phi)_+$
    emit accepted draft tokens and the corrected token; discard the remaining draft; start a new round
  end if
end for
if all K drafts were accepted, emit them and sample one extra token from the target distribution

The draft acceptance rate α_{acc} determines how much proposed work survives. Under a simplified model in which each drafted position is accepted independently with that probability, a round emits its accepted prefix plus one target-sampled token, either the correction at the first rejection or the extra token after full acceptance. Each draft position contributes a token only if every earlier draft was accepted, so each successive position adds one more factor of α_{acc} and the expectation is a geometric series over the $K+1$ possible emissions. Equation 10.6 gives its closed form, the expected output of one verification round with K drafted tokens:

$$\mathbb{E}[N_{\text{emit}}] = \frac{1-\alpha_{\text{acc}}^{K+1}}{1-\alpha_{\text{acc}}} \tag{10.6}$$

The target streams its weights once per round whatever the acceptance, so the weight traffic amortized per emitted token falls in proportion to equation 10.6. Low acceptance can make generation slower than ordinary decoding, because rejected drafts still cost drafting and verification time. That mechanism belongs to one replica. A serving fleet must decide when to enable speculation, how its memory cost interacts with admission, and whether it improves the endpoint's SLO-compliant capacity.

Those decisions start from how steeply the benefit falls with acceptance. The following calculation assumes $K = 5$ draft tokens, a draft step $20\times$ faster than a target decode step, and a verification pass costing one target decode step, which isolates sensitivity to acceptance. A round then costs K draft steps plus one verification step, and the speedup divides the expected tokens emitted per round by that round cost, with the round cost measured in target decode steps:

High acceptance ($\alpha_{\text{acc}} = 0.9$) yields an expected 4.69 tokens per verification round for a 3.7× speedup. Medium acceptance ($\alpha_{\text{acc}} = 0.7$) delivers 2.94 tokens per round for a 2.4× speedup. Low acceptance ($\alpha_{\text{acc}} = 0.5$) drops progress to 1.97 tokens per round, yielding a modest speedup of 1.6×.

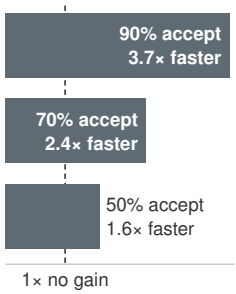

Speedup shrinks quickly as draft acceptance falls.

Acceptance depends on the draft-target pair, sampling policy, and request distribution. The preceding calculations are sensitivity cases, not fixed acceptance rates for chat, reasoning, or code generation. A serving fleet must measure acceptance and useful token throughput for its actual traffic before selecting a speculation policy, and that policy must answer three serving concerns: memory, the SLO, and change over time.

The first concern is memory. The draft model is not free state. It requires memory for weights, and each in-flight request that uses it reserves additional draft KV state. Pairing a 7B draft model with a 70B target adds roughly 14 GB of FP16 weights per replica plus per-request draft cache. When colocated with the target, this memory overhead competes directly with the KV-cache admission budget developed in section 10.7.1; every gigabyte reserved for the draft is unavailable to target requests. Under a binding memory limit, that reservation lowers maximum concurrency. Placing the draft on separate devices preserves target memory at the cost of additional hardware and communication. The admission controller and the speculation policy therefore require a joint capacity and latency evaluation.

The second concern is which service-level objective the endpoint must meet. At sufficient acceptance, speculation can reduce TPOT, while draft execution and cache residency reduce admission capacity, so the scheduler should compare SLO-compliant goodput with and without speculation under the intended load. TTFT needs separate measurement because draft setup and prefill can add work before the first response token.

The third concern is that acceptance and load change over time. Batch size moves the payoff as well as acceptance does. Verification is cheap only while decode leaves compute idle, and a larger batch consumes that headroom, so the same draft configuration can cross from a gain to a loss as concurrency rises. Variable accepted lengths also add scheduler work. The crossover batch must be measured for each draft-target pair and machine rather than inferred from a batch-one result. A scheduler can enable speculation when measured TPOT or goodput improves and disable it when verification cost outweighs accepted progress. Disabling draft execution, however, frees neither the resident draft weights nor the draft KV pages. Recovering that capacity requires unloading the weights and ending or migrating the speculative state. Self-speculative decoding, Medusa prediction heads (Cai et al. 2024), and lookahead decoding change this budget, and each needs its own quality, acceptance, and resource evaluation.

10.7.7 Preemption and the victim policy

When active sequences grow their contexts until every free KV-cache page is claimed, the replica can neither advance those sequences nor admit a higher-priority arrival without evicting a sequence that is already running. Iteration-level admission removes the batch-barrier form of head-of-line blocking, because finished sequences leave and waiting ones join at every iteration boundary. Memory pressure can still produce head-of-line blocking under strict first-come-first-served (FCFS) order, because a waiting request too large for the free KV pages holds back smaller requests that would fit. Once the KV pool fills, admitting a waiting request or extending an active one means taking capacity from a running sequence.

Preemptive scheduling allows the system to stop allocating decode iterations to one request so that another can run. Preemption is practical at token boundaries, but it is not free. The scheduler must retain the request's KV pages in HBM, swap them to host memory, or discard them and repeat prefill when the request resumes. A useful victim policy compares those three costs with the waiting request's remaining deadline rather than preempting at a fixed generation length. Retaining state consumes scarce capacity, swapping consumes interconnect bandwidth, and recomputation consumes prefill compute; under heavy load, rejecting low-priority work can cost less than repeatedly moving or rebuilding its state.

10.7.8 KV-cache budgeting and eviction

Admitting sequences up to a naive average-length estimate without reserving headroom for context growth forces the replica into constant mid-generation preemptions or out-of-memory crashes. Effective KV-cache management therefore begins with capacity accounting rather than an eviction policy. The serving process first reserves memory for weights, activations, runtime overhead, and safety headroom, then treats the remaining HBM as a managed pool for active sequences:

```
Available GPU memory = Total - Weights - Activations - Overhead
KV pool size = 0.9 * Available  # Leave 10% headroom

Max concurrent sequences = KV pool size / (avg_seq_length * KV bytes per token)
```

That budget estimates how many sequences can coexist, but an average length does not reserve space for future growth. Admission must account for the admitted requests' remaining budgets. LRU eviction can reclaim retained prefix blocks that no active request references. Reclaiming an active sequence's state instead requires a preemption decision, made with the victim policy of section 10.7.7. What capacity accounting adds is the observation that freeing the most bytes is not necessarily the cheapest recovery, because the largest sequence is also the most expensive to swap or rebuild.

Example 10.4: KV cache memory hierarchy

Scenario: A multi-tenant LLM serving system experiences KV-cache HBM exhaustion (80 GB limit) during peak concurrency.

Diagnosis: Pure HBM allocation limits node concurrency to 50 active sequences in this scenario. Implementing a 3-tier memory hierarchy (HBM → CPU DRAM → Non-Volatile Memory Express (NVMe) SSD) allows swapping inactive sequence pages to CPU DRAM or NVMe. The illustrative swap budgets of 1–5 ms and 10–50 ms, respectively, require validation for the bytes transferred, effective bandwidth, and contention; they are not fixed properties of the storage tiers.

Systems lesson: In this scenario, multi-tier paging raises retained sequence capacity from 50 to 500, while the active HBM-resident set remains limited to 50. Suspended sequences must be restored before execution resumes under this policy. The extra retention does not establish a throughput gain; restoration costs must fit the affected requests' deadlines. Streaming active state from another tier would be a different execution policy with transfer costs on the decode path.

10.7.9 Adapter state and locality

Deploying a separate copy of a 140 GB base model for each of 10,000 fine-tuned tenants is prohibitive in HBM, yet fetching a tenant's weight delta from host memory on the critical path stalls every decode step. Alongside the KV-cache attention state of concurrent sequences, adapter serving therefore adds a second claim on the replica's memory and bandwidth budget. A low-rank adaptation (LoRA) adapter is a user- or task-specific weight delta that must be available when its request executes. It stores that delta as the product of two thin matrices, so it holds far fewer values than the weight matrix it modifies. LoRA modifies a shared base model without copying all of its weights: the serving process keeps the 140 GB base resident once while loading the adapters needed by active requests. Whether adapter storage or transfer becomes limiting depends on the active adapter mix, ranks, and reuse.

> **Systems Perspective 10.2: The context switch of machine learning**
>
> Adapter locality becomes a performance constraint when a scheduled request needs weights that are not resident. Loading those weights from host memory resembles part of an operating-system context switch, because execution must wait for the required state unless transfer overlaps other useful work. An adapter already in HBM need not be transferred from the host on every generation step. Coordinating admission with adapter placement can reduce these stalls, while grouping requests by adapter can improve reuse at the cost of waiting for compatible work. The policy must balance locality against the latency and fairness of less frequent adapters.

10.7.9.1 Multi-LoRA serving and batched adapter execution

In multi-tenant environments, thousands of fine-tuned adapters share a single base foundation model. For a linear projection layer with base weights $\mathbf{W}_0 \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$, an adapter k parameterizes the forward transformation using low-rank decomposition matrices $\mathbf{A}_k \in \mathbb{R}^{r_k \times d_{\text{in}}}$ and $\mathbf{B}_k \in \mathbb{R}^{d_{\text{out}} \times r_k}$ with rank $r_k \ll \min(d_{\text{in}}, d_{\text{out}})$ (typically $r \in \{8, 16, 64\}$), as formulated in equation 10.7:

$$\mathbf{y} = \mathbf{x}\mathbf{W}_0^T + \frac{\gamma_k}{r_k}\mathbf{x}\mathbf{A}_k^T\mathbf{B}_k^T \tag{10.7}$$

where $\mathbf{x} \in \mathbb{R}^{1 \times d_{\text{in}}}$ is a row vector, $\mathbf{y} \in \mathbb{R}^{1 \times d_{\text{out}}}$, and γ_k is a fixed scaling factor, the hyperparameter usually written α in LoRA papers and renamed here because α already denotes startup latency and draft acceptance in this chapter. Each adapted projection adds $r_k(d_{\text{in}} + d_{\text{out}})$ parameters instead of a full $d_{\text{in}}d_{\text{out}}$ weight matrix. The resident adapter budget therefore depends on which layers are adapted, their dimensions and ranks, precision, and the number of simultaneously loaded tenants.

Packing requests with different adapters into one batch lets the base projection $\mathbf{X}\mathbf{W}_0^T$ execute as one shared dense general matrix multiply (GEMM), but a naive implementation then launches a separate small kernel for each request's adapter $(\mathbf{A}, \mathbf{B})$. Hundreds of such launches per layer make host dispatch and launch latency significant even when each adapter performs little arithmetic, and the small kernels leave GPU streaming multiprocessors (SMs) underused. The following batched operation removes the need for one launch per request.

↪ How low-rank adapters are trained, and how their rank is chosen under a device's memory and energy budget, is developed in section 11.4.3.2.

> **Definition 10.4: Batched gather-matrix-vector multiplication (BGMV)**
>
> ***Batched gather-matrix-vector multiplication (BGMV)*** is a GPU kernel operation that batches matrix-vector products across distinct adapters, dynamically gathering adapter weights from a memory pool according to a per-token adapter index vector.
>
> 1. **Significance**: It reduces per-request kernel launch overhead in multi-adapter serving, allowing different adapters to share kernel launches. The remaining overhead depends on adapter ranks, batch composition, and memory locality.

2. **Distinction**: A single dense GEMM applies one weight matrix across its input rows. BGMV gathers the adapter-specific weight slice for each element, combining distinct matrix-vector products in one kernel operation.
3. **Common pitfall**: Batched execution does not itself solve allocation waste. Loading and unloading differently sized adapters can fragment a contiguous allocator. Paged adapter storage addresses that separate problem, provided the execution kernels support its layout.

Systems such as Punica (Chen et al. 2024) and S-LoRA (Sheng et al. 2024) address launch overhead and adapter memory through two architectural mechanisms (definition 10.4):

Batched gather-matrix-vector (BGMV) kernels Instead of launching separate kernels for every request, batched or segmented gather-matrix-vector operations group different adapters within each projection. Down-projection and up-projection can use separate kernels; Punica distinguishes shrink and expand operators (Chen et al. 2024). The scheduler builds a per-request adapter index, and each kernel uses it to gather a request's activation row and its adapter's matrices, apply the rank-r down- and up-projections of equation 10.7, and add the resulting delta to the shared base projection's output. The gain must be measured against the actual adapter mix; irregular access and small ranks can still limit achieved bandwidth.

Unified paged adapter memory management S-LoRA (Sheng et al. 2024) jointly manages adapter weights and KV state in a unified GPU memory pool. Its page holds a vector along the shared hidden dimension, rather than a fixed operating-system-sized byte allocation. Adapter matrices $\mathbf{A}_k$ and $\mathbf{B}_k$ occupy noncontiguous pages, allowing adapters of varying ranks ($r \in \{8, 16, 32, 64\}$) to share capacity with growing KV caches. Paging reduces fragmentation; both kinds of state still consume the same finite pool.

The published design keeps adapters in host memory and uses the waiting queue to predict which weights the next batch will need. Prefetching overlaps adapter transfers with current execution when sufficient time and free GPU memory are available. A cache miss whose transfer cannot be hidden still delays the affected request.

Across a tensor-parallel group, S-LoRA aligns each adapter with the base layer's column- or row-parallel split (section 6.5.3) and communicates only the r_k-dimensional intermediates rather than full d_{out}-dimensional activations. The layout avoids replicating adapter weights across GPUs at the price of one extra collective on the small rank-r_k intermediate.

Which collective that is depends on the base layer's split. When the base layer is column-parallel, the input $\mathbf{x}$ is already replicated across GPUs, so S-LoRA partitions $\mathbf{A}_k$ across the rank dimension, uses an AllGather to assemble the full r_k-dimensional intermediate on every GPU, and partitions $\mathbf{B}_k$ across the output columns to match the base weights. When the base layer is row-parallel, the input $\mathbf{x}$ arrives partitioned across the input dimension, so S-LoRA splits $\mathbf{A}_k$ along that same input dimension, uses an AllReduce to sum the partial r_k-dimensional vectors (Sheng et al. 2024), and applies $\mathbf{B}_k$ before the final output assembly shares the base layer's own collective.

KV pages and adapter deltas share HBM with the base weights, so every byte freed in one allocation can admit work in another. When a request claims that memory, and how fast it then uses it, depends on the phase it is in: prefill builds a request's cache in one pass, while decode extends it one token at a time.

10.8 Prefill and Decode

Admitting a 4,096-token prompt onto a worker that is already streaming replies to thirty users saturates the GPU's Tensor Cores for hundreds of milliseconds and freezes every active token stream until the prompt finishes. A language-model request is two workloads in sequence. Prefill reads the whole prompt in one forward pass and writes its keys and values into the cache sized in section 10.7.1, while decode extends that cache one token per step until the response ends. Because the two phases run the same weights on opposite sides of the roofline ridge, their asymmetry decides which optimization helps which phase, how a replica schedules them when they share a worker, and whether they should share one at all.

10.8.1 The phase asymmetry

That roofline split comes directly from how many tokens each phase feeds through a linear projection at once. During the prefill phase, the engine computes key-value representations across all prompt tokens (S_{prompt}) concurrently. Because all prompt tokens are processed in parallel, the linear projection GEMM operations operate on large sequence dimensions ($M = S_{\text{prompt}}$). For a projection layer with hidden dimension d, arithmetic intensity scales with prompt length:

$$I_{\text{prefill}} \approx \frac{2 \cdot S_{\text{prompt}} \cdot d^2}{2 \cdot d^2} = S_{\text{prompt}} \quad \left[\frac{\text{FLOPs}}{\text{byte of weights loaded}}\right]$$

A 1,024-token prompt therefore performs about 1,024 FLOPs per byte of FP16 weights, well past an H100's FP16 ridge. Under this weight-traffic-only approximation, increasing S_{prompt} beyond the accelerator's ridge point moves the projection's roofline ceiling to the compute limit. Reaching that ceiling still depends on kernel efficiency; short prompts, attention, other memory traffic, and communication can limit the complete prefill phase, so prefill as a whole still requires profiling.

In contrast, the subsequent decode phase generates a single token per sequence sequentially, so the GEMM's row dimension is the number of sequences decoding together, not the prompt length. At batch one the intensity is about 1 FLOP per byte of FP16 weights. At the decode batch sizes this chapter considers, from tens of sequences up to the knee near 80 in notebook 10.2, the arithmetic intensity $I_{\text{decode}} \approx B_{\text{dec}}$ stays below the ridge point. This places the weight projections in a memory-bandwidth-bound regime under this approximation. Each decode iteration streams the active weights from HBM and reuses them across the batch, so the cost of each generated token falls as the batch grows until KV-cache capacity stops further admission, which is why batch size is the decode phase's control knob. The approximation omits KV-cache reads, which grow with every sequence's context, so larger batches and longer contexts add bytes to every step and hold decode on the bandwidth slope short of the ridge.

↩ The ridge point that separates compute-bound from bandwidth-bound execution is established in section 2.3.2.

The asymmetry also sorts the optimizations by phase. Tiled, IO-aware attention reduces attention's memory traffic mainly during prefill, because only prefill forms the full prompt-by-prompt score matrix that a tiled kernel avoids writing to HBM (section A.5.2). A decode step has a single query row and no quadratic intermediate; it benefits instead from a fused kernel that reads each cached key and value once and computes scores, softmax, and the weighted sum in one pass. Speculative decoding (section 10.7.6) targets decode alone, and KV-cache quantization reduces decode's storage and read traffic while its encoding cost begins when prefill writes the cache. Precision helps both phases through different mechanisms. A weight-activation format raises the compute ceiling that long-prompt prefill reaches, while low-bit weights shrink the bytes that weight-dominated decode streams (section 10.9). Each gain depends on format and kernel support, conversion overhead, and the costs the change leaves untouched, and each must preserve output quality.

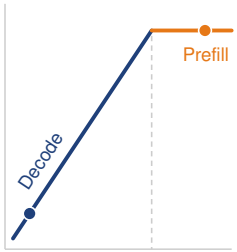

Long-prompt prefill reaches the compute ceiling; decode stays on the bandwidth slope.

These are properties of each phase in isolation. Once continuous batching admits new prompts onto a worker that is decoding other streams, the asymmetry becomes a scheduling conflict between them, and a replica has two ways to resolve it: interleave the phases on one worker, or give each phase its own pool.

10.8.2 Chunked prefill and prompt-decode interleaving

Continuous batching admits work at iteration boundaries, yet production requests arrive with widely disparate prompt lengths. Admitting a long prompt in one piece holds the worker for its entire compute-bound prefill while the bandwidth-bound decode streams beside it wait.

> Definition 10.5: Chunked prefill
>
> ***Chunked prefill*** is an LLM serving mechanism that divides long prompt evaluations into discrete token chunks of size at most C_{chunk}, interleaving prefill computation with ongoing decode iterations within a shared token budget.
>
> 1. **Significance**: It limits the prompt-decode interference problem where monolithic prefills stall active decode streams for hundreds of milliseconds, giving the scheduler control

over how much prefill work can lengthen the TBT of active streams. Meeting interactive streaming service-level objectives (SLOs) still requires measuring the resulting tail latency.

2. **Distinction**: Unlike continuous batching with monolithic prefill (which prefills each prompt in one uninterrupted pass before the worker returns to decoding), chunked prefill treats prefill as a multi-step iterative process co-scheduled alongside memory-bound decode tokens.
3. **Common pitfall**: Sizing the chunk budget C_{chunk} too small incurs kernel launch and attention mask overheads, degrading overall prefill throughput. Sizing it too large allows prefill compute to monopolize execution time, causing TBT tail spikes.

With an unchunked prefill policy, the scheduler runs an incoming prompt through the whole model before returning to decode work. The interference comes from scheduling that whole operation without interleaving other requests' decode steps, and the delay lands in the TBT of requests already decoding on the same worker. A long response can dilute an isolated stall in its average TPOT even though the user sees a conspicuous gap, so prioritizing the new prompt's time to first token can consume the streaming-latency budget of requests already decoding. The measured prefill duration and the endpoint's TBT target determine whether interleaving is necessary.

10.8.2.1 The chunked prefill scheduling mechanism

To limit this interference, chunked prefill (definition 10.5, used by Sarathi-Serve (Agrawal et al. 2024)) partitions prompts into chunks of at most C_{chunk} tokens. A candidate range such as 256–512 tokens must be tested against the model, hardware, and streaming deadline rather than treated as a universal setting. A prompt of length S_{prompt} requires at least $\lceil S_{\text{prompt}}/C_{\text{chunk}} \rceil$ scheduling iterations; a smaller residual token budget can require more.

The two schedules in figure 10.13 hand the same long prompt to the same worker. Run whole, its prefill sits between two decode steps of every active stream, so each stream's token gap stretches to the full prefill time. Split into chunks, the prompt advances one chunk per hybrid iteration beside the streams' decode tokens, so each stream waits at most one hybrid iteration, while the prompt's own first token arrives later than it would have under the monolithic schedule.

In each iteration, the scheduler constructs a hybrid batch combining active decode requests with prefill chunks under a global iteration token budget N_{budget}, a count of tokens rather than a time, as formulated in equation 10.8:

$$B_{\text{dec}} + \sum_{i \in \mathcal{P}_{\text{active}}} C_i \leq N_{\text{budget}} \tag{10.8}$$

where B_{dec} is the number of active decode tokens and C_i is the number of prompt tokens allocated to waiting prefill request i, so every term counts tokens. The simplified loop assumes admission control reserves sufficient KV-cache capacity and a future decode slot for each admitted request, keeping the decode count within the iteration budget even after prefills are promoted. Each waiting request contributes at most one chunk per iteration, with its committed prompt offset unchanged until execution completes. Algorithm 10.3 formalizes this iteration-level hybrid scheduling loop.

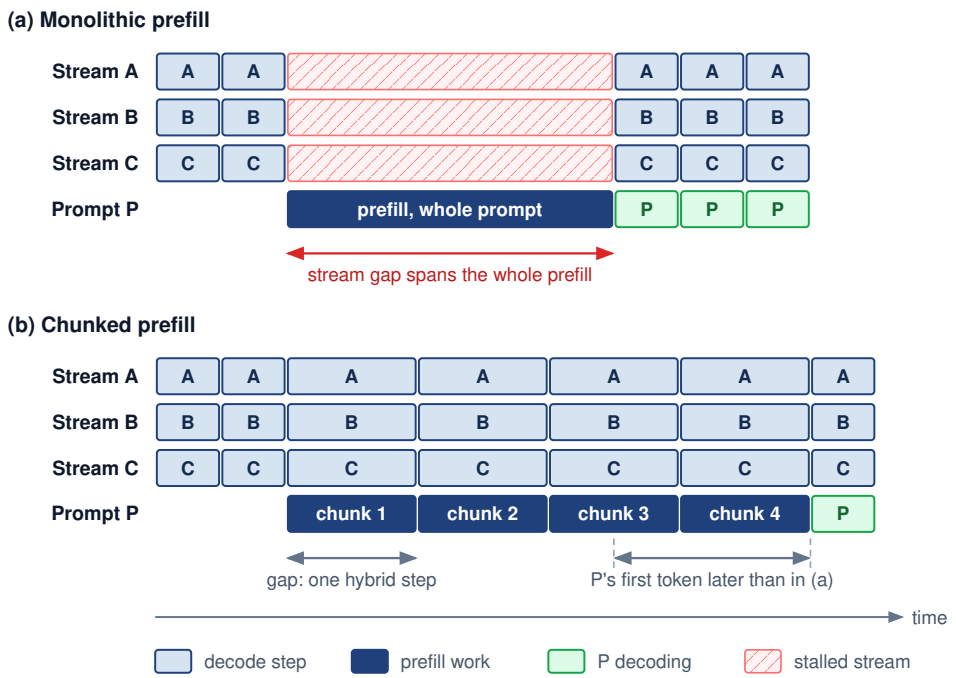

Figure 10.13: Chunked Prefill Trades One Prompt's TTFT for Every Stream's TBT: Three streams decode while a long prompt P arrives; the timing is illustrative. (a) A monolithic prefill runs P's whole prompt in one pass, and every active stream stalls for its full duration. (b) Chunked prefill gives P one chunk per hybrid iteration beside the streams' decode tokens, so each stream's gap shrinks to one hybrid iteration, and P's first token arrives later. Chunking changes when the prefill work runs, not how much of it there is.

Algorithm 10.3 Chunked prefill and hybrid batch scheduler

Require: waiting prefill queue $\mathcal{Q}_{\text{prefill}}$; active decode set $\mathcal{D}$; iteration token budget N_{budget}; maximum chunk size C_{chunk}
Ensure: generated tokens for decode requests; incremental KV-cache construction for prefills

while serving engine is active **do**
  $B_{\text{dec}} \leftarrow |\mathcal{D}|$
  $N_{\text{rem}} \leftarrow N_{\text{budget}} - B_{\text{dec}}$
  $\mathcal{B}_{\text{prefill}} \leftarrow \emptyset$
  for each R_i in a snapshot of $\mathcal{Q}_{\text{prefill}}$, in queue order **do**
    if $N_{\text{rem}} > 0$ **then**
      $o_i \leftarrow R_i$.offset
      $C_i \leftarrow \min(C_{\text{chunk}}, N_{\text{rem}}, R_i.\text{prompt_length} - o_i)$
      $\mathcal{B}_{\text{prefill}} \leftarrow \mathcal{B}_{\text{prefill}} \cup \{(R_i, o_i, C_i)\}$
      $N_{\text{rem}} \leftarrow N_{\text{rem}} - C_i$
    end if
  end for
  run hybrid forward pass using each prefill slice $R_i[o_i : o_i + C_i]$ and the decode set $\mathcal{D}$
  append newly generated tokens for all $R \in \mathcal{D}$
  remove completed decode requests from $\mathcal{D}$ and release their KV blocks
  for $(R_i, o_i, C_i) \in \mathcal{B}_{\text{prefill}}$ **do**
    update prompt offset: R_i.offset $\leftarrow o_i + C_i$
    if R_i.offset $= R_i$.prompt_length **then**
      promote R_i from $\mathcal{Q}_{\text{prefill}}$ to $\mathcal{D}$
    end if
  end for
end while

10.8.2.2 Weight-loading amortization via decode piggybacking

Executing prefill and decode tokens together in shared projection operations amortizes weight loading across both phases, an effect termed **decode piggybacking**. Alternating separate prefill and decode operations limits scheduling interference but does not by itself establish this weight reuse.

In a small-batch isolated decode step, streaming the model parameters (M_{weights} bytes) from HBM to registers achieves low arithmetic intensity because only B_{dec} tokens share the weight reads. Memory bandwidth then limits how much work the compute units can sustain, leaving arithmetic capacity available for a larger shared projection.

In a hybrid iteration, the model weights loaded from HBM into on-chip cache and registers are simultaneously reused for both the compute-dense prefill chunks ($\sum_i C_i$ tokens in total) and the memory-bound decode tokens (B_{dec} tokens). Ignoring attention-state traffic, communication, and scheduling overhead gives the projection-only lower-bound estimate in equation 10.9:

$$T_{\text{step}} \gtrsim \max\left(\frac{M_{\text{weights}}}{\text{BW}_{\text{mem}}}, \frac{2 \cdot (B_{\text{dec}} + \sum_i C_i) \cdot P_{\text{model}}}{R_{\text{peak}}}\right) \tag{10.9}$$

where P_{model} is the active parameter count, BW_{mem} is HBM memory bandwidth, and R_{peak} is peak accelerator compute throughput. These quantities must describe the same device or shard. Balancing the total prefill token allocation against weight streaming can amortize weight-loading cost across both phases, but this lower bound cannot guarantee $T_{\text{step}} \leq \text{SLO}_{\text{TBT}}$. Chunk sizing must leave room for the omitted costs and be validated against measured tail latency under the intended workload.

Example 10.5: Sarathi: Chunked prefill implementation

Scenario: An LLM inference server receives a mix of long input prompts ($10{,}000+$ tokens) and interactive short decode requests (Agrawal et al. 2024).

Diagnosis: When the worker schedules a full long-prompt prefill before returning to decode, that prefill can occupy execution for hundreds of milliseconds and create a TBT spike for active streams.

Systems lesson: Chunking moves the cost of a long prompt rather than removing it. At 512 tokens per chunk, a $10{,}000$-token prompt needs at least $\lceil 10{,}000/512 \rceil = 20$ iterations to finish its prefill. With this prompt as the only prefill request, active streams share each decode iteration with at most one of its chunks, while the prompt's own time to first token stretches across those 20 iterations. Multiple waiting prompts can contribute additional chunks within the total iteration budget. The chunk budget sets that exchange between other streams' TBT and the long prompt's TTFT, so it must be tuned against both measured tails.

10.8.3 Prefill-decode disaggregation

Even with chunked prefill, every hybrid step still forces active decode streams to wait while prompt chunks execute on the same GPU, and a single worker pool must compromise between prefill's compute-heavy parallelism and decode's bandwidth- and capacity-heavy layout. Disaggregation removes that compromise by giving each phase its own hardware pool, while the KV-cache sizing of section 10.7.1 sets the price of separating them. The cache that prefill builds is the exact payload that the network must hand off to the decode pool.

Disaggregation partitions the serving pipeline into specialized physical tiers (figure 10.14). In the illustrated long-prompt, small-decode-batch regime, the compute-bound prefill pool and the bandwidth-bound decode pool are joined by a remote direct memory access (RDMA) KV-cache transfer (Zhong et al. 2024). This keeps the phases from competing on one accelerator and permits independent parallelism and capacity plans. The phase asymmetry tells each pool what to buy. The prefill pool wants arithmetic throughput for its compute-bound GEMMs, while the decode pool wants memory bandwidth and capacity for weight streaming and KV cache, which also makes it the natural home for aggressive weight quantization. The transfer and queueing at each pool still add to request latency, and other batch or context sizes can change which resources bind.

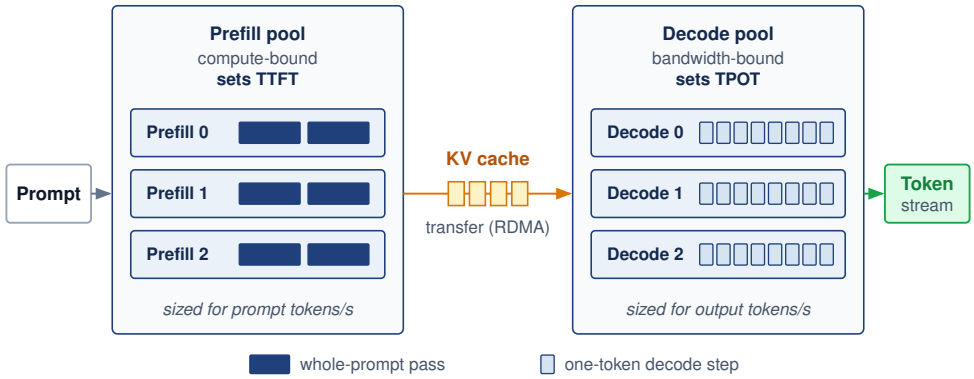

Figure 10.14: Disaggregated Serving Architecture: A long-prompt, small-decode-batch regime separates compute-bound prefill instances from memory-bandwidth-bound decode instances, mediated by RDMA KV-cache transfer. Because prefill sets time to first token and decode sets time per output token, each pool is sized for its own token rate. Other batch and context sizes can change the binding resources.

The disaggregation decision therefore depends on a quantitative contract, not on the phase labels alone. If prefill produces M_{KV} bytes of attention state and the usable inter-pool bandwidth is β, then the handoff consumes at least

$$T_{\text{handoff}} = \alpha + \frac{M_{\text{KV}}}{\beta}, \tag{10.10}$$

where α includes transfer setup and queueing for the operating point being evaluated; it is not a fixed network latency once load changes. The payload term sets the scale. A 128K-token request on the GQA model of notebook 10.5 carries tens of gigabytes of KV state, so moving it over a single 450 GB/s NVLink direction takes on the order of a hundred milliseconds, and an inter-node link takes longer still. Spreading the transfer across a tensor-parallel group divides that time only when each device has its own path to its counterpart in the decode pool. Because M_{KV} grows linearly with prompt length, the long prompts that make disaggregation attractive are also the prompts whose handoff costs most, and that cost lands on TTFT.

Equation 10.10 therefore makes placement part of the service contract. The prefill pool must sustain the admitted prompt-token rate, the decode pool must sustain the admitted output-token rate, and both pools need headroom for the request-length distribution. Disaggregation justifies its KV-cache transfer cost only when eliminating phase interference and balancing the two pools raises SLO goodput. Otherwise, colocated chunked prefill avoids the additional handoff and pool coordination.

Whether the phases share a worker or split into pools, the base weights remain the largest fixed claim on each device's memory and, for decode, the largest share of the bytes streamed per step. Weight quantization addresses that part of the budget, with numerical quality and execution cost determining how much of the storage saving becomes useful serving capacity.

10.9 Weight Quantization for Serving

Keeping a 70B model's 140 GB of FP16 weights intact forces every replica across at least two 80 GB GPUs and makes every decode step wait on streaming 140 GB from HBM. Weight quantization for serving is therefore a bottleneck-selection decision: reducing precision turns that 140 GB FP16 weight payload into a 35 GB INT4 payload, allowing a model that previously required multiple GPUs to fit on one (perspective 10.3). The ideal payload reduction is 4×, although metadata, runtime buffers, and KV state still enter the deployment's memory budget. Because neither quality loss nor speedup is fixed by nominal bit width, both must be measured for the deployed representation and workload.

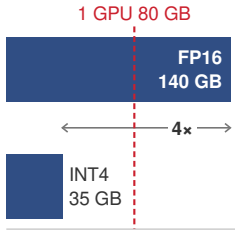

INT4 turns a two-GPU model into a one-GPU candidate.

Systems Perspective 10.3: Quantization as a serving constraint

In serving systems, quantization changes the binding resource as well as the model's size. Post-training quantization (PTQ) can shrink a deployed model without retraining, while quantization-aware training (QAT) spends training effort to preserve accuracy at lower precision. QAT simulates low-precision rounding in the forward pass and passes gradients through the rounding step as if it were absent, the straight-through estimator (Bengio et al. 2013), so the weights adapt to the error they will carry. It earns that cost only after a PTQ recipe misses the quality budget at the target precision, and a recipe that passes at INT8 says nothing about INT4. The practical order is to deploy with PTQ when it meets quality and reach for QAT, or adapter-based fine-tuning of a quantized base, when it does not, rechecking quality after the final deployment conversion. The systems question is which constraint the lower-precision representation relaxes: GPU memory capacity, HBM bandwidth during decode, accelerator instruction support, or cost per served token.

Choosing a quantized representation starts with deciding which quantity should shrink: resident weights for fit, bytes moved per decode step for bandwidth, activations for prefill throughput, or the precision path supported by target hardware. Serving at scale also introduces distinct operational constraints: a deployment team may need post-training quantization without the original training corpus, quality must be preserved across diverse inputs, and hardware deployment targets vary from data center GPUs to edge accelerators. Lack of the original corpus does not imply data-free calibration; the methods that follow may still require representative deployment inputs. Production inference therefore starts from a decision model rather than from the methods. Table 10.11 ties each method family to the serving constraint it relaxes; what changes for each when it is deployed is its calibration cost, kernel path, and the evaluation a serving team must run. Method names such as W4A16 and W8A8 give the bit widths of the weights (W) and of the activations (A).

Table 10.11: Serving Quantization Decision Model: Quantization methods are useful only when they relax the resource that binds the serving system. The method family follows from the bottleneck: fit, bandwidth, prefill compute, or outlier robustness.

Binding serving constraint	Representation change	Method family	Risk to test before rollout
Model does not fit in memory	Quantize weights only	GPTQ, AWQ	Calibration mismatch and quality regression
Decode is bandwidth-bound	Reduce bytes read per token	W4A16 weight-only paths	Kernel support and KV-cache headroom
Prefill is compute-bound	Quantize weights and activations	SmoothQuant W8A8	Activation outliers and INT8 hardware path
Outliers block low precision	Change the coordinate basis before quantization	Rotation-based methods	Transform overhead and limited runtime path

10.9.1 Outliers and block-wise scaling

When a single weight or activation channel is ten times larger than its neighbors, uniform per-tensor quantization must either clip that extreme value or stretch its step size across the entire range, collapsing the surrounding small values to zero. Large language models produce such outlier features routinely: in the models studied by Dettmers et al. (2022), about 0.1 percent of hidden dimensions carried activation magnitudes roughly 3–20$\times$ larger than the rest. Weight tensors face the same damage when a few large weights set the scale for the whole matrix. Block-wise quantization contains that weight-side damage inside the stored format.

Definition 10.6: Block-wise quantization

Block-wise quantization is an ML quantization scheme that partitions a weight tensor into nonoverlapping groups of b_{block} elements and assigns each group its own scale. An affine scheme uses $s_i = (x_{\max,i} - x_{\min,i})/(2^b - 1)$ together with an offset or zero point; a symmetric

scheme centers the range at zero and needs only scale metadata. The metadata examples below assume symmetric scale-only quantization.

1. **Significance**: With block size $b_{\text{block}} = 64$ and $b = 4$ bits, weights compress from 16 bits to 4 bits (a $4\times$ memory reduction) while each FP16 scale adds $16/64 = 0.25$ bits per weight. This yields an effective bit-width of 4.25 bits per weight: 6.25 percent overhead relative to the INT4 payload, or about 1.6 percent of the original FP16 weight size.
2. **Distinction**: Unlike per-tensor quantization, which applies a single scale across the entire weight matrix, block-wise quantization contains outlier damage within individual blocks. Per-tensor scaling forces that scale to accommodate outlier values, which wastes precision on the majority of near-zero weights. A single extreme value therefore cannot degrade quantization fidelity for the whole tensor.
3. **Common pitfall**: A frequent misconception is that block size is a free hyperparameter. Smaller blocks ($b_{\text{block}} = 32$) can improve local range resolution but double the scale metadata overhead vs. $b_{\text{block}} = 64$. At $b_{\text{block}} = 16$, the scale adds 1 bit per weight: 25 percent overhead relative to the INT4 payload, or 6.25 percent of the original FP16 weight size. The effect on model quality and kernel speed must be evaluated rather than inferred from block size alone.

The metadata price stays small across formats. An INT8 scheme with one FP32 scale per 128 weights adds $32/128 = 0.25$ bits per weight, about 3 percent of its payload. Block-wise scaling therefore isolates outliers without the bandwidth penalty of a higher-precision format, which keeps the decode traffic saving of a low-bit format intact. Figure 10.15 isolates this dynamic-range protection across an eight-weight sample containing a single outlier (58.7). Under per-tensor scaling (left), the outlier dictates a coarse scale factor ($s \approx 0.462$) that forces six of the seven small weights to round to zero. Partitioning the tensor into four-element blocks (right) confines that coarse scale to the second block, allowing the uncorrupted first block to adopt a fine scale ($s_1 \approx 0.00165$) that preserves full INT8 resolution.

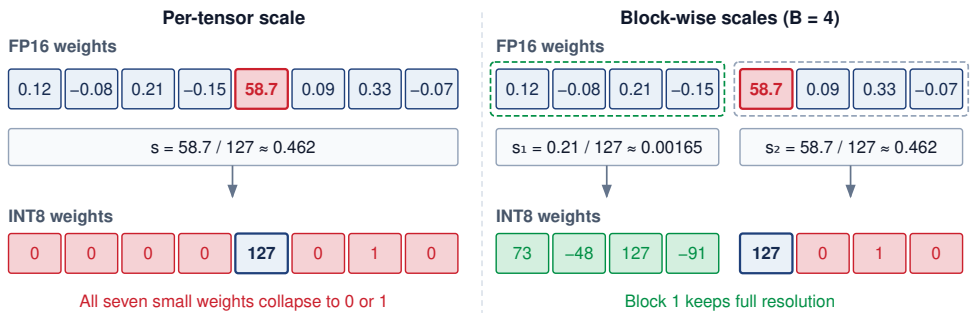

Figure 10.15: Block-Wise versus Per-Tensor Quantization: Under per-tensor scaling (left), one outlier weight sets the scale for the whole tensor, so six of the seven small weights round to zero and the seventh rounds to one. Under block-wise quantization (right), each group of weights gets its own scale: the block without the outlier keeps full INT8 resolution, while the small weights that share the outlier's block are still crushed. Smaller blocks confine the damage further at the cost of more scale metadata.

10.9.2 LLM-specific quantization challenges

Block-wise scaling handles weight outliers inside the stored format. The outlier activation problem is harder, because activations arrive at run time and their extreme channels cannot be rescaled offline block by block. A coarse scale loses resolution for smaller values, while a narrow range clips large ones. This trade-off is not unique to language models, but its channel structure motivates the methods below.

For a serving deployment the consequence is that no single INT8 scale suits both the outlier channels and the typical values. The method families therefore differ in which part of the serving

contract they protect, at which stage they pay for it, and what they cost to deploy. LLM.int8() (Dettmers et al. 2022) pays at run time. It detects outlier dimensions against a magnitude threshold, multiplies them in FP16 and the rest in INT8, and sums the two products. This preserves quality at the cost of outlier detection and a second execution path on every step.

GPTQ (Frantar et al. 2023) and AWQ (Lin et al. 2024) pay offline and both produce 4-bit weight-only artifacts. GPTQ quantizes a layer's weights in sequence and adjusts the still-unquantized weights to compensate for the error already introduced, guided by a second-order estimate from calibration activations of how sensitive the layer's output is to each weight. AWQ instead protects the salient weights, those that multiply large activations, because a small rounding error on a weight channel that multiplies a large activation produces a large output error. Multiplying those salient weight channels by a scale factor greater than one before group quantization forces them to span a larger portion of the shared 4-bit integer grid, reducing their relative rounding error while an inverse scale is applied to the corresponding input activations. Both methods need representative calibration inputs and supported packing kernels, a reconstruction objective that is only a proxy for task quality, and a comparison at the same model, bit budget, and held-out evaluation. AWQ also needs its inverse scales fused into the preceding layer of the deployed graph, or it spends part of the latency saving on an extra runtime operation.

SmoothQuant (Xiao et al. 2023) targets low-precision matrix multiplication rather than weight-only storage and pays in a preprocessing transform. It divides each activation channel by a smoothing factor and multiplies the matching weights by the same factor. This leaves the product unchanged before rounding and migrates part of the extreme activation range into the weights, so that both tensors can fit inside INT8 for W8A8 execution. That matters for compute-bound prefill, although shrinking the activation range stretches the weight range and makes the weights harder to quantize. The smoothing factor therefore balances the two quantization errors, and production activation extremes can still differ from the calibration ranges. Rotation-based methods such as QuaRot (Ashkboos et al. 2024) apply orthogonal Hadamard transforms that spread large coordinates across channels, which makes W4A4 a serving candidate at the price of some online rotation work.

None of these properties fixes quality. Each artifact must be compared with the FP16 baseline on the same model, evaluation data, and runtime, reporting its calibration set, group size, kernel path, and any online-transform overhead. Its serving benefit must also be measured as SLO-compliant throughput on the target workload rather than inferred from nominal bit width.

10.9.3 W8A8 versus weight-only ceilings

Quantizing only the weights (W8A16 or W4A16) cuts weight traffic during decode but leaves compute-bound prefill executing on FP16 Tensor Cores at unchanged arithmetic speed, whereas quantizing both weights and activations (W8A8) risks accuracy loss from activation outliers without further shrinking weight bytes. The prefill-decode split therefore decides which path can pay: both schemes shrink the weights, but only W8A8 changes the arithmetic ceiling. Table 10.12 bounds what each 8-bit path can buy relative to the FP16 baseline.

Table 10.12: W8A8 SmoothQuant vs. Weight-Only Quantization: Ceilings derived from bytes per weight and the H100's INT8-to-FP16 peak ratio, not measured speedups. Both schemes halve weight bytes, so both can at most halve weight traffic in bandwidth-bound decode. Only W8A8 moves prefill onto INT8 Tensor Cores. Realized gains and quality loss depend on kernels, calibration, and the request mix.

Configuration	Weight bytes	Prefill compute ceiling	Decode weight-traffic ceiling	Quality
FP16 (baseline)	1×	1×	1×	Baseline
W8A16 (weights only)	Reduced 2×	1× (FP16 math)	Up to 2×	Measure vs. FP16
W8A8 (SmoothQuant)	Reduced 2×	Up to 2× (INT8 Tensor Cores)	Up to 2×	Measure vs. FP16

The distinction is which ceiling binds. Among these configurations, W8A8 enables INT8 matrix multiplication, with the illustrated 2× peak ratio providing a compute ceiling rather than a measured

prefill speedup. Both quantized schemes halve the weight payload. If weight reads dominate decode, activation quantization does not further reduce that traffic, although it can change other memory and compute costs. Quantization and dequantization also add work. The end-to-end benefit therefore depends on the prompt/output mix, batch sizes, and measured execution path.

Hardware and software support decide which of these savings reach the serving workload, because the arithmetic format and the resident storage format need not be identical. A common INT4 weight-only kernel reads packed weights and dequantizes them to FP16 on chip during the GEMM, so it cuts weight capacity and read traffic without any INT4 arithmetic. For batch-one decode dominated by weight reads, throughput then scales roughly as memory bandwidth divided by the model's weight bytes, omitting KV-cache traffic and execution overhead, and leaving the arithmetic in FP16 costs little because arithmetic is not the binding term. At large batch or in prefill, where execution approaches the compute ceiling, the gain comes instead from a weight-activation format, W8A8 or FP8 on hardware with FP8 Tensor Cores, which raises the ceiling that weight-only formats cannot move. In the ideal memory-bound case, 4-bit weights cut weight traffic by 4×, but realized decode throughput depends on kernel support, dequantization overhead, KV-cache precision, batch shape, and whether HBM bandwidth is truly the bottleneck. A serving system can therefore run different formats at different operating points, with each configuration evaluated against the same quality criterion and its own latency targets, and the value of 4-bit serving is strongest when it increases model residency or batch capacity.

10.9.4 Runtime contract for quantized serving

Loading a 4-bit checkpoint into a serving runtime that expands the weights back into resident FP16 tensors at startup or falls back to unoptimized kernels forfeits both the memory capacity and the HBM bandwidth savings before a single request arrives. Production serving frameworks determine whether a quantized artifact retains its intended advantages by converting packed weights inside the execution kernel rather than in HBM. The serving stack must therefore report the resident representation and kernel path, budget KV state around actual allocations, and expose quality and latency regressions (table 10.13).

Table 10.13: Quantized Serving Runtime Contract: Framework-specific APIs differ, but every production path has to preserve representation, choose supported kernels, plan memory, and validate quality under the deployed traffic mix.

Runtime responsibility	Why it matters
Preserve the quantized artifact	Full-model expansion into resident FP16 forfeits the reduced weight storage and traffic.
Select compatible low-precision kernels	Unsupported operators fall back to slower or higher-precision execution.
Plan memory around weights and KV cache	The capacity gain matters only if the freed memory becomes useful batch headroom.
Validate quality and latency together	A lower-bit path that meets latency but fails task quality is not a serving win.

Weight quantization reduces the weight allocation, while PagedAttention reduces waste in allocating KV state. Both feed the same admission budget as the capacity estimator of notebook 10.5, with quantized weights in place of FP16 ones. That estimate is optimistic, because a usable integer limit must round down after reserving runtime buffers, quantization metadata, and safety headroom, and sharded deployments must satisfy the budget on every device.

This comparison uses a conservative decimal HBM budget matching the A100 nameplate label, rather than the registered byte capacity used by the node estimator. A 70B model with 4-bit weights requires approximately 35 GB, leaving 45 GB on an 80 GB A100 for KV cache. With FP16 KV cache at about 2.6 GB per 1K-token sequence for the MHA baseline, this supports about 17 concurrent 1K-token sequences before other runtime overheads. With FP16 weights, the same 70B model would not fit on one A100 at all, so weight-only quantization changes both feasibility and batch capacity.

10.9.5 Selecting a quantization format

Choosing a quantization format from datasheet TFLOPS or nominal bit width alone fails when a 4-bit checkpoint degrades task accuracy on production prompts or when a smaller replica completes

fewer SLO-compliant tokens per GPU. The final selection must satisfy memory, quality, and latency requirements together. Native support makes a precision path a candidate without establishing the model's quality at that precision, and no universal quality ordering separates 4-bit weight-only from 8-bit weight-and-activation quantization, so both are evaluated against the same task baseline. For compute-bound prefill, W8A8 offers the illustrated 2× INT8 compute ceiling; for decode dominated by weight reads, W4A16 can cut the weight payload by up to 4×, provided conversion and other costs leave a useful saving.

These precision choices affect cost per million useful tokens. The A100 residency calculation explains why fewer GPUs become possible at all, because INT4 weights shrink the resident payload enough for a 70B replica to fit on fewer devices. Table 10.14 isolates the effect of GPU count by assuming equal per-GPU prices and the same completed token rate across three configurations. Under that contract, reducing the number of billed GPUs reduces cost proportionally. The calculation does not establish that the smaller replicas sustain the same rate or meet the same quality and latency targets.

Table 10.14: Quantization Impact on Serving Cost: Estimated cost per million tokens at three configurations. At equal unit prices and unchanged useful token throughput, halving the GPU count halves the cost; halving it again brings the per-million-token cost to one quarter of the FP16 baseline. The cost savings come almost entirely from fewer GPUs rather than from per-GPU throughput gains.

Configuration	Cost per 1M tokens (estimated)
FP16 on 8× A100	$2.40
AWQ 4-bit on 4× A100	$1.20
AWQ 4-bit on 2× A100	$0.60

The illustrated 2–4× cost reduction follows from the assumed device-count reduction at unchanged useful throughput. If the smaller replica completes fewer qualifying tokens per second, cost per token falls by less and can increase. A deployment comparison must therefore combine billed replica cost with measured SLO-compliant output under the same quality contract.

Precision moves the capacity wall without removing it. A 175-billion-parameter model still needs about 87.5 GB for its 4-bit weight payload alone, more than one 80 GB H100 holds before KV state and runtime allocations. Keeping that representation accelerator-resident requires distributing it across devices. Offloading to host memory is another placement choice, but introduces transfers that must fit the latency budget. Sharding addresses the capacity problem while adding communication and coordination to each replica.

10.10 Model Sharding for Inference

When a deployment must retain a 70B model's 140 GB of FP16 weights in accelerator memory, or when a 175B model's 87.5 GB INT4 weight payload still exceeds a single 80 GB GPU, no single accelerator can hold the replica. Where quantization shrinks the representation, sharding keeps it intact and divides it across multiple GPUs that execute together as one logical replica. The resulting design adapts tensor and pipeline parallelism from training, but must evaluate their communication and scheduling costs against per-request inference deadlines.

10.10.1 Memory and latency triggers for model sharding

A single-GPU replica fails when model weights overflow its HBM capacity or when its serial forward pass takes too long to meet the endpoint's TTFT and TPOT targets. Model sharding addresses the memory and latency rows of table 10.1 through different mechanisms, and the timing analysis below determines whether parallel execution repays its communication cost.

The first driver is memory capacity. If the selected representation must remain in accelerator memory but does not fit on one device, the replica needs multiple devices. For a model with P parameters at precision b bits, equation 10.11 gives the weight payload before metadata and runtime state:

$$M_{\text{weights}} = P \times \frac{b}{8} \text{ bytes} \tag{10.11}$$

Applied to the 70B model in FP16 (16 bits), equation 10.11 recovers the payload quoted since section 10.2. Each parameter occupies two bytes, so the unquantized checkpoint alone requires:

$$M_{\text{weights}} = 7 \times 10^{10} \times \frac{16}{8} = 140 \text{ GB}$$

The result exceeds the 80 GB capacity of an H100 GPU, establishing a weight-only lower bound of 2 shards for an accelerator-resident replica. KV state, buffers, replication within the chosen partition, and uneven placement can require a larger degree.

The second driver is latency. Even when a model fits in memory, sharding can reduce latency by parallelizing computation. Equation 10.12 is the Fleet Law (equation 1.3) applied to a single inference request, separating the compute term that sharding divides from the communication and synchronization terms it adds:

$$T_{\text{step}}(N) = \frac{T_{\text{compute}}}{N} + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}} \tag{10.12}$$

where N is the number of devices in the sharding group, $T_{\text{comm}}(N)$ is the communication overhead, $T_{\text{sync}}(N)$ is synchronization overhead, and T_{overlap} is communication hidden behind useful compute. In the common no-overlap simplification with negligible extra synchronization, this reduces to $T_{\text{compute}}/N + T_{\text{comm}}(N)$. Sharding provides latency benefit only when the nonoverlapped overhead is smaller than the time saved through parallelization.

10.10.2 Tensor parallelism

Splitting every linear projection independently across GPUs would force an AllReduce synchronization after every matrix multiplication, stalling decode on interconnect latency. Tensor parallelism (Shoeybi et al. 2019) avoids half of those collectives by splitting each layer with the column-row partitioning developed for training. This keeps the activation between paired linear layers local and costs two AllReduce operations per transformer block, one after attention and one after the feed-forward sublayer.

The attention computation partitions across heads, so with N_{heads} heads and tensor-parallel degree N_{TP} each device computes attention for $N_{\text{heads}}/N_{\text{TP}}$ of them when the head count divides evenly. The KV cache is sharded with the heads, so each device stores cache only for its own heads. A multi-head attention layout spreads many KV heads across the group, while grouped-query attention with eight KV heads leaves one per device in an eight-way group.

↩ Column-row tensor partitioning and its two AllReduce operations per transformer block are established in section 6.5.3.

```
8-way tensor parallelism, 64 KV heads (MHA):
Device 0: KV cache for heads 0-7
Device 1: KV cache for heads 8-15
...
Device 7: KV cache for heads 56-63
```

Because each device holds cache only for its own heads, reusing that cache on other devices, as a handoff from a prefill pool to a decode pool does, requires compatible cache content and placement. Matching tensor-parallel configurations can align the head partitions, but the receiver must also use the same model state, prefix inputs, cache representation, and logical token positions. A different partition requires redistribution rather than direct shard reuse. Equal parallelism degree alone does not establish cache validity.

Each device applies its shard of the attention output projection to its local heads, and an AllReduce sums the projected partial outputs; its latency includes collective startup and coordination as well as payload transfer. Figure 10.16 traces the resulting two synchronization points per layer across two GPU ranks.

The inference time with tensor parallelism follows equation 10.13:

$$T_{\text{inference}} = \frac{T_{\text{compute}}}{N_{\text{TP}}} + 2 \times T_{\text{allreduce}}(M) \tag{10.13}$$

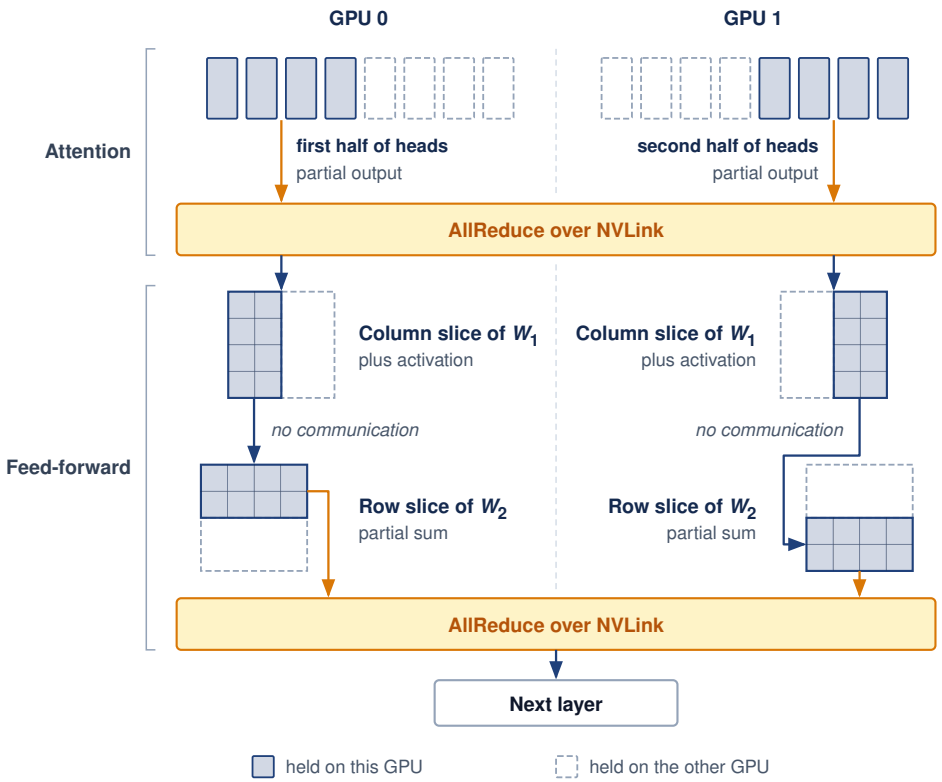

Figure 10.16: Tensor Parallelism for Inference: Computation is distributed across devices by splitting tensor operations. Attention heads are partitioned across GPUs, requiring an AllReduce operation to synchronize results. Feed-forward networks use a column-row splitting strategy that requires only one AllReduce synchronization per block. This partition can reduce latency when saved compute exceeds communication cost, which is why the two collectives run over a fast local interconnect.

Equation 10.13 is the per-layer instance of equation 10.12. Here T_{compute} is now the sequential compute time of one layer, and M is the full activation payload each rank contributes to each AllReduce, priced by the ring cost model of section 5.6.3. Compute divides ideally across ranks, but the reduced activation keeps its full shape, (tokens in the step) $\times$ (hidden width) $\times$ (bytes per element), and the factor of 2 counts the attention and feed-forward collectives (example 10.6). When a model spans multiple nodes without high-bandwidth NVLink across every rank, those two per-layer collectives become prohibitively expensive, motivating inter-layer pipeline partitioning instead.

↩ Hiding tensor-parallel collectives behind compute on one node, and what that overlap can and cannot recover, is established in section 9.7.2.

Example 10.6: Tensor parallelism for Llama-70B

Scenario: An engineering team deploys Llama-70B (70B parameters, FP16 footprint 140 GB) across an 8-GPU H100 node with NVLink (Touvron et al. 2023).

Diagnosis: The calculation assumes 30 ms of serial layer compute as a reference, independent of the devices required to hold the weights: 12 ms for attention and 18 ms for the feed-forward block. That reference describes a long-prompt prefill, in which every prompt token passes through the layer in one step, so it is far larger than one layer's share of a decode step on the reference replica, which passes only one new token per sequence through the layer. Ideal eight-way division gives 1.5 ms and 2.25 ms, and the two AllReduce operations of equation 10.13 add an assumed 0.3 ms each, for 4.35 ms per layer.

Systems lesson: The illustrative eight-way model gives a 6.9× reduction relative to its serial-compute reference (prefill time reduced from 2,400 ms to 348 ms). The entire shortfall from the ideal 8× is the two collectives, a fixed cost per layer that dividing compute cannot shrink. Their share of the layer therefore grows as per-device compute falls, which is why a decode step, with far less compute per layer than this prefill reference, feels the collective cost most.

10.10.3 Pipeline parallelism for inference

Splitting an 80-layer model sequentially across four GPU stages leaves three accelerators idle during an isolated request and does nothing to shorten its forward pass. Pipeline parallelism distributes layers across devices so each stage runs a contiguous subset of layers and passes activations to the next stage without intra-layer synchronization.

For inference, pipeline parallelism creates bubbles differently than in training. Maintaining multiple in-flight requests fills the stages that an isolated request leaves idle and amortizes the initial bubble across steady-state token delivery (figure 10.17).

↩ The pipeline bubble, the time stages sit idle while a pipeline fills and drains, and the microbatching that shrinks it in training are established in section 6.5.2.

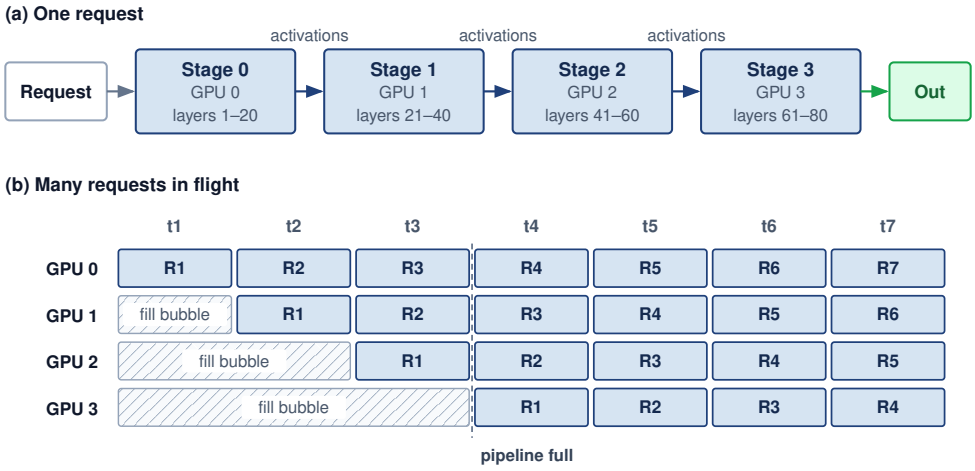

Figure 10.17: Pipeline Parallelism Bubbles: (a) An 80-layer model split into four stages passes one request's activations from stage to stage, so splitting the layers alone does not shorten its sequential computation; it must still traverse all four stages. (b) With many requests in flight, later stages idle during the fill bubble, and from slot t4 onward every stage is busy and one request leaves the pipeline per slot. Pipeline parallelism therefore suits throughput-oriented serving better than latency-critical interactive serving.

For a single request, pipeline staging does not parallelize the dependent layers. If stage execution times are unchanged, the request still pays their sum plus transfers, so staging alone provides no execution-latency reduction. A comparison against a memory-constrained implementation that offloads weights can differ because keeping stages resident changes those execution costs.

Pipeline parallelism instead scales throughput by sustaining multiple concurrent requests across execution stages, as tracked in the steady-state schedule of figure 10.17(b). Once the pipeline reaches saturation, throughput is limited by the slowest stage. For N_{PP} balanced stages with negligible transfer overhead, this is N_{PP} times the throughput of executing the whole model sequentially on one device, not N_{PP} times a single stage's throughput. An individual request still traverses every stage, so its execution latency remains approximately the sum of the stage times.

Pipeline parallelism therefore suits inference when memory forces sharding but latency targets are relaxed, when aggregate throughput matters more than individual request latency, or when the interconnect handles point-to-point transfers better than fast AllReduce. The two sharding approaches differ in single-request latency, throughput, communication pattern, activation memory, and implementation complexity (table 10.15).

Table 10.15: Pipeline vs. Tensor Parallelism: Key trade-offs between layer-internal tensor splitting (can reduce single-request latency when compute savings exceed collective cost) and inter-layer pipeline staging (overlaps concurrent requests but introduces pipeline bubbles). Throughput is relative to sequential execution of the whole model on one device; per-layer AllReduces keep tensor parallelism below the $N_{TP}\times$ ideal, and balanced, full pipelines approach $N_{PP}\times$.

Aspect	Tensor parallelism	Pipeline parallelism
Single-request latency	Compute term reduced up to $N_{TP}\times$	No reduction from overlapping this request alone
Throughput	Below $N_{TP}\times$ (AllReduce cost)	Up to $N_{PP}\times$ (pipelined)
Communication pattern	AllReduce (bandwidth-intensive)	Point-to-point (latency-sensitive)
Memory efficiency	Activations replicated	Activations passed along
Complexity	Higher (requires custom kernels)	Lower (layer-level partitioning)

The decision starts with capacity feasibility and then compares latency and throughput on the available topology. Neither a fast interconnect nor a relaxed deadline alone establishes the best choice; the request mix and measured execution costs determine the useful partition.

10.10.4 Expert parallelism for MoE models

An MoE model activates only a small fraction of its parameters per token, yet replicating the entire expert pool on every GPU quickly exhausts device memory. Mixture-of-experts (MoE) models such as DeepSeek-V3 (DeepSeek-AI 2025) and Mixtral (A. Q. Jiang et al. 2024) replace each feed-forward network with parallel expert sub-networks and a lightweight router that selects a few experts per token, growing parameter capacity without proportionally growing expert arithmetic per token. Routing, shared layers, and communication still add work. Expert parallelism places different experts on different GPUs to distribute their storage and execution when the model and workload exceed a single device.

10.10.4.1 MoE economics and capacity planning

MoE serving separates two capacity measures. The *total parameter count* sets the weight-storage requirement at the chosen precision, while the *active parameter count* estimates the weight-related arithmetic and traffic for each token. Neither count includes all runtime costs. An accelerator-resident deployment must accommodate inactive experts as well as active ones, even though each token uses only a subset.

The following comparison isolates weight-related costs at batch size 1. Streaming all weights of a dense 400B FP16 model moves 800 GB per step. Counting only DeepSeek-V3's 37B active parameters gives approximately 74 GB in FP16, a 10.8× reduction in weight bytes per token. The proportional estimate of weight-related arithmetic is 10.8× lower. These are payload and arithmetic comparisons, not equivalent-quality model comparisons or measured speedups; cache traffic, expert placement, routing, and communication affect realized latency.

The trade-off is resident capacity: keeping all experts on accelerators avoids fetching inactive experts when routing changes. Offloading can reduce that resident footprint but adds transfer and scheduling costs. DeepSeek-V3's full model in FP16 requires approximately 1342 GB, requiring many devices if the entire FP16 weight set remains accelerator-resident.

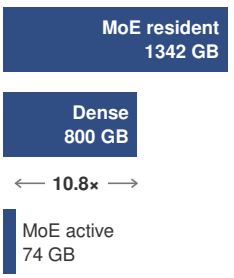

MoE keeps every expert resident but activates few per token.

Napkin Math 10.8: MoE capacity planning

This comparison uses a conservative decimal budget of 80 GB per GPU, matching the H100 nameplate label rather than the registered byte capacity used by the node estimator.

Problem: A team deploys DeepSeek-V3 (671B total parameters, 37B active per token) for a chatbot application. The model uses FP8 weights (1 byte per parameter). The cluster has 8-GPU nodes, each with 8× H100 GPUs (80 GB HBM per GPU, 640 GB per node). How many nodes are needed, and what is the per-token decode latency?

Math:

1. **Memory**: Total weight memory = 671 GB. A single node (640 GB) is insufficient. 2 nodes provide 1,280 GB in aggregate. The remaining capacity must still accommodate KV state, metadata, buffers, and any replicated weights on each device.
2. **Latency**: Count 37B active parameters (37 GB) per decode step. If their reads were evenly distributed across 16 GPUs, each would read about 2.3 GB. At the stated 3.35 TB/s HBM bandwidth, $T_{\text{read}} \approx 0.69$ ms. The scenario adds 0.3 ms for routing. Expert placement and token routing do not guarantee this balance, especially at batch size 1.
3. **Floor**: The simplified sum is ≈ 1.0 ms/token, corresponding to 1,000 tokens/s at batch size 1 under these assumptions. It is an optimistic estimate, not a measured rate or a rigorous latency bound; compute, KV traffic, synchronization, and imbalance remain unmodeled.

Systems insight: Sparse activation separates resident parameter capacity from the work selected for one token. Planning an accelerator-resident 671B model therefore requires both a per-device memory budget and a timing model for the actual expert placement. Dividing total active bytes by device count does not establish that every device can contribute bandwidth equally.

10.10.4.2 Expert parallelism and load balancing

Expert parallelism carries over from training unchanged in mechanism. Each GPU holds a subset of experts, and each token travels to the devices holding the experts its gate selects. What changes in serving is who pays for imbalance, because an overloaded expert now delays live requests rather than a training step. In an MoE layer, a gating network selects k_{experts} of the N_{experts} experts for each token and sums their outputs weighted by its gate probabilities. An even expert placement assigns each device N_{experts}/N experts, where N is the number of expert-parallel devices, and each token then makes the round trip of figure 6.19. One AllToAll dispatches its representation to the ranks hosting its experts, and a second returns their outputs for the weighted sum.

The communication pattern differs from tensor parallelism. AllReduce delivers one combined result to every device, so per-link traffic is fixed by the message size. Expert parallelism uses AllToAll to deliver different routed tokens to different devices, so link and receiver load depend on the router's choices, and the host of a popular expert becomes an incast point whose completion time the other hosts wait on (section 5.8.1). Balancing the number of stored experts does not balance token traffic.

Serving inherits training's balance controls without their guarantees. The auxiliary balancing loss shapes routing on training data, not on future traffic. The expert capacity factor caps accepted tokens per expert (1.25× is an illustrative setting), leaving the server to decide what happens to the excess. Dropping or rerouting it changes the model computation and needs quality validation, while expert buffering retains it at the cost of latency and queue capacity.

10.10.4.3 Routing failure modes

Router behavior affects both system performance and model quality. Expert collapse during training concentrates traffic on a few experts, and at serving time a shifted input distribution can create hotspots without any change to the stored model. Token dropping under a capacity limit makes the computation differ from the uncapped model, so its quality impact must be evaluated together with any throughput gain. Because routing selects experts dynamically,[5] the routing pattern becomes a systems workload rather than only a model-internal choice.

↩ The capacity factor, the auxiliary load-balancing loss, and the hot-expert skew they bound during training are established in section 6.5.5.

[5] **MoE (mixture-of-experts) serving trade-off**: Only a subset of experts activates per token (2 of 8 in Mixtral, 8 of 256 routed experts in DeepSeek-V3). Keeping all weights accelerator-resident gives a total-to-active weight ratio of roughly 3.6× (Mixtral) to 18× (DeepSeek-V3). Expert parallelism distributes that storage while adding token-routing communication; offloading is an alternative that trades resident capacity for weight transfers.

Example 10.7: Expert parallelism for Mixtral-8x7B

Scenario: Mixtral-8x7B (46.7B total parameters, ~12.9B active per token) is served across a 4-GPU group using top-2 routing.

Diagnosis: Replicating all 8 experts on every GPU requires the full 93.4 GB FP16 weight set per device, more than one 80 GB H100 holds. Placing 2 experts on each GPU (4-way expert parallelism) cuts each device's share to about 25.8 GB, the shared attention layers plus its local

experts. Every MoE layer, however, now needs an AllToAll dispatch and gather to move tokens to the devices that hold their selected experts.

Systems lesson: Expert parallelism trades AllToAll communication for per-GPU parameter memory. The trade pays only when the collective, repeated at every MoE layer, fits inside the decode budget that the smaller resident footprint makes possible.

Expert parallelism couples state placement to input-dependent access. Embedding-heavy recommendation models face a related problem: entity popularity determines which stored vectors receive traffic. Their sharding strategy must therefore balance storage and lookup load, just as expert placement must account for both resident weights and routed tokens.

10.10.5 Embedding sharding for recommendation systems

Production recommendation models carry multi-terabyte embedding tables that cannot fit on any single server and would stall every ranking query if fetched over thousands of serial network hops. Meta's DLRM reference architecture applies model parallelism to those tables while scaling dense layers separately (Naumov et al. 2019), and serving inherits the row-wise, column-wise, and hybrid layouts that training uses to shard them (section 6.5.6).

What changes at serving time is the constraint the layout answers to. Training shards tables for sparse update throughput, while serving must finish every lookup within a per-request deadline. Popularity skew concentrates lookups on the shards that own hot keys, and under a latency SLO the slowest of those shards sets the tail of every request that touches it. Serving layouts therefore replicate or column-partition hot entries even at extra storage cost, keep the cold tail row-sharded, and batch lookups by destination shard so that one request does not wait on thousands of sequential round trips.

The following scenario uses Archetype B (DLRM at Scale), a recommendation model built around large embedding tables, and combines row-wise, column-wise, and hybrid sharding. A single Archetype B model's embedding tables occupy 10 TB or more. The platform-wide resident budget of 100 TB is deliberately set well above that, so that the sharding arithmetic is stressed at the upper end of the regime.

Lighthouse 10.1: Archetype B (DLRM at Scale): Embedding sharding budget

Consider an Archetype B recommendation workload with a large entity catalog and a smaller resident embedding tier:

Scale:

- Resident embedding budget (platform-wide): 100 TB
- Entity catalog: 10 trillion identifiers
- Embedding dimension: 128–256
- Shards: 1,000+ servers

The arithmetic of these bullets is the design constraint. A dense INT8 table with 10 trillion entities and 128 dimensions would require at least 1,280 TB, while 100 TB across the same entity count averages only 10 bytes per entity, a fraction of one 128-byte vector. Most entities therefore cannot hold a resident vector at all. The resident tier keeps the frequently accessed subset, and the long tail lives in another tier or a compressed representation, which is exactly the split the sharding strategy below is built around.

Sharding strategy:

The tier percentages below describe the admitted resident population, not the full entity catalog. Replica bytes count against the platform-wide resident budget, so the admitted population and replication scope must be chosen together.

- **Hot embeddings** (top 1 percent by access frequency): Replicated within selected shard groups

- **Warm embeddings** (next 10 percent): Column-sharded with 8-way parallelism
- **Cold embeddings** (remaining 89 percent): Row-sharded with consistent hashing

Each inference request requires approximately 5,000 embedding lookups. Instead of one RPC per lookup, the batched design accumulates lookups for 1 ms, deduplicates entities across requests, groups lookups by destination shard, dispatches to all shards in parallel, and assembles embeddings as responses arrive.

Illustrative before-and-after assumptions for sequential dispatch versus parallel shard batching appear in table 10.16.

Table 10.16: Illustrative Embedding Lookup Comparison: Assumed latency and bandwidth compare sequential per-lookup RPCs with shard-batched parallel dispatch. One dispatch wave still contains multiple network requests; it does not mean one RPC to the entire fleet.

Metric	Without optimization	With optimization
Sequential dispatch waves	5,000 individual RPCs	One parallel wave; multiple shard RPCs
Lookup latency	50 ms	2 ms
Network bandwidth	10 Gbps	40 Gbps (burst)

The lighthouse lesson is that sparse embedding placement and lookup coordination can dominate DLRM-scale serving even when dense arithmetic has spare capacity. The illustrated workload motivates measuring shard traffic and lookup latency before adding dense compute.

10.10.6 Hybrid sharding strategies

Extending tensor parallelism across node boundaries stalls decode on slow inter-node AllReduces, while using pipeline parallelism alone leaves single-request latency untouched. Production systems therefore combine sharding strategies into hybrid sharding, using one axis to cut per-token latency inside a node, another to fit weights across nodes, and a third to route sparse components without forcing dense layers onto the wrong device. Joint placement across these axes is established in section 6.6, so the serving question is which axis buys latency. Tensor and pipeline parallelism combine when a model needs both memory distribution and latency reduction:

```
8 GPUs organized as 2 by 4 (pipeline stages by tensor parallel):

Stage 0 (Layers 1-40):  TP across GPUs 0,1,2,3
Stage 1 (Layers 41-80): TP across GPUs 4,5,6,7
```

The four-way tensor split can cut per-layer compute time by up to $4\times$, while the two pipeline stages complete the eight-way memory distribution; pipelining adds no per-request latency benefit of its own. Expert-parallel MoE layers and recommendation models that keep embeddings on CPU servers beside a tensor-parallel dense model follow the same rule, and the recommendation split reappears in the case study of section 10.15.2.

Hybrid sharding buys fit or latency by adding new communication patterns at every boundary between shards. The next question is whether that added communication time stays inside the latency and throughput budget.

10.10.7 Interconnect latency and communication overhead

Sharding a model across more GPUs can make inference slower rather than faster when collective synchronization takes longer than the compute time it saves. Each sharding strategy has a characteristic communication pattern, and each pattern carries a cost model developed earlier in this book. Pipeline parallelism passes activations point to point between stages and pays one startup latency plus a bandwidth term per transfer (section 3.4.4). Tensor parallelism pays a ring AllReduce per collective, whose startup term grows linearly with the number of devices while its bandwidth term approaches twice the payload divided by link bandwidth (section 5.6.3). Expert parallelism pays an

AllToAll, whose completion time is set by the largest of per-message startup, endpoint injection or reception, and traffic across the fabric bisection (section 5.8).

These models give an initial communication budget, not a complete prediction of request latency. Skewed expert traffic, network contention, or a different collective algorithm can change the realized timing, so the chosen topology and collective implementation must sustain that budget under the intended concurrent load.

Communication overhead depends heavily on the interconnect technology, because each candidate link fixes the physical bandwidth, latency, and scope available to the sharded model (table 10.17).

Table 10.17: Data-Center Interconnect Specifications: Bandwidth, latency, and intended scope for NVLink, PCIe Gen5, InfiniBand HDR, and 100G Ethernet, establishing the physical limits that bound achievable sharding speedups.

Interconnect	Bandwidth	Latency	Use case
NVLink (H100)	900 GB/s	500 ns	Intra-node TP
PCIe Gen5	64 GB/s	1 µs	Intra-node (no NVLink)
InfiniBand HDR	200 Gb/s (25 GB/s)	7 µs	Inter-node
Ethernet 100G	100 Gb/s (12.5 GB/s)	50 µs	Inter-node (commodity)

For tensor parallelism the latency column, not the bandwidth column, decides placement. Decode divides a small per-layer compute time across the group while still paying two collectives per layer for every generated token, so the fixed per-collective latency in table 10.17 sets the overhead share. The 0.3 ms per collective assumed in example 10.6 already lies an order of magnitude above the bandwidth-only transfer time on NVLink, so startup, synchronization, and launch costs dominate it. An inter-node link with higher per-message latency would multiply that cost on every token. That decode regime is why tensor-parallel groups are normally kept inside one NVLink domain.

10.10.8 Selecting a sharding strategy

Choosing the wrong sharding axis either burns interconnect bandwidth on per-layer collectives that the link cannot support or leaves single-request latency above the SLO. No approach wins on latency priority, throughput priority, interconnect limits, and implementation effort at once (table 10.18).

Table 10.18: Sharding Strategy Selection Guide: Evaluation of tensor, pipeline, expert, and embedding sharding across latency targets, throughput priorities, interconnect constraints, and implementation complexity.

Factor	Tensor parallel	Pipeline parallel	Expert parallel	Embedding shard
Latency priority	Compute savings exceed collectives	No staging-only reduction	Routing-dependent	Lookup-dependent
Throughput priority	Collective-limited	Full, balanced pipeline	Balanced expert traffic	Balanced lookup load
Interconnect limited	Poor fit	Good fit	Moderate	Good fit
Implementation effort	High	Low	Moderate	High

Strategy selection starts from the deployment bottleneck. Tensor parallelism is a candidate when reduced layer execution can repay frequent collectives; pipeline parallelism distributes layers and can sustain throughput when enough concurrent work fills the stages. Expert and embedding sharding distribute sparse state while adding routed communication. These choices can be combined, but each must fit the per-device memory budget and the measured latency contract.

Once a replica's memory footprint and execution capacity are established, demand determines how many replicas the service needs. The next decision is where to send each arriving request. Routing must account for work already in progress and reusable state, so that replica capacity becomes useful fleet capacity.

10.11 Load Balancing and Request Routing

Suppose ten model replicas are serving traffic when a request for a 5,000-word document summary arrives. Assigning it to an overloaded replica adds waiting and can increase interference for work already admitted there. The effect depends on admission, batching, and priority policies. Round-robin assignment does not inspect those differences in remaining work. A state-aware router uses memory and queue information to choose a feasible destination, while limiting the cost and staleness of its observations.

Lighthouse 10.2: Archetype B (DLRM at Scale): The tail at scale

Archetype B (DLRM at Scale) uses p99 $\leq$ 100 ms as its illustrative latency target. At 10^7 QPS, an event probability of 10^{-4} implies an expected 1,000 such events per second. This count motivates examining the tail, but does not identify its cause. Two-choice routing can reduce imbalanced queuing delay $L_{\text{lat,queue}}$; traces must establish whether queueing, inference time $L_{\text{lat,infer}}$, or another component drives the misses.

Load balancing serves two goals that can conflict: sending each request to the replica that can serve it fastest, and spreading load so that no replica idles or saturates. Latency-optimal routing can concentrate load on fast replicas until they slow down. Evaluation therefore needs signals for both goals, queue length and estimated remaining work per replica for latency and utilization spread for balance. It also needs the decision overhead of the routing choice. All signals are judged against completed-request tail latency and admission outcomes rather than balanced request counts.

10.11.1 Tail latency under fan-out

When a serving request fans out to 100 backend shards that each meet their own p99 latency target independently, the assembled response misses that threshold about 63 percent of the time. Many serving requests are not one call but a set of parallel sub-requests whose results must all return before the response can be assembled. Examples are a ranking query that scores candidates on many shards, a recommendation that looks up embeddings across tables, and a sharded model whose replicas must all finish a step. Because the request completes only when its slowest leaf returns, it inherits the tail of every leaf it touches, the way a series system fails when any one component fails. If each of N_{series} leaves independently exceeds its own latency threshold with probability p, the request exceeds it with probability $1-(1-p)^{N_{\text{series}}}$.

Reading the same expression backward turns it into a budgeting rule. To keep the whole request within its p99 target, each of the 100 leaves may be slow only about 0.01 percent of the time, so the leaf target moves from p99 to roughly p99.99. Fan-out width is therefore a tail-latency cost that the service pays per request, and the levers act on different terms. Fewer, larger shards lower N_{series}. Queue-aware routing and isolation lower p by keeping leaves out of their own queueing tails. A deadline with a partial or default result, as in the search cascade of section 10.15.4, caps the damage when a leaf still misses. Hedged requests lower p further at the price of duplicated work. The independence assumption is optimistic, because a shared switch, host, or hot key slows many leaves together, so the measured composed tail is the quantity to budget, not the formula alone.

10.11.2 Round-robin and random assignment

Routing requests in a fixed cycle or at random works when every call costs the same, but in an inference fleet where one request generates 10 tokens and the next generates 2,000, stateless assignment piles long queues onto unlucky replicas while others sit idle. Round-robin cycles through the servers in order, which distributes work evenly only when servers are homogeneous and processing times identical, while random assignment picks a server uniformly, converging to an even mean with higher variance. Both are reasonable baselines for identical servers, but heterogeneous GPU generations, variable request sizes, warmup behavior, and replicas nearing memory limits make the state of the server fleet matter in production.

A static placement model isolates the imbalance caused by random choice: place R equal-size jobs sequentially into R initially empty bins, without departures. With uniform independent choices,

↩ The same independence model applied to storage reads, and hedged reads with their failure-domain placement, are established in section 4.3.6.

the expected maximum bin occupancy follows the classical balls-into-bins result (Mitzenmacher 2001) in equation 10.14:

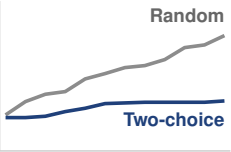

Two-choice routing flattens the tail imbalance of random placement.

$$\mathbb{E}[\text{max bin occupancy}] = \Theta\left(\frac{\log R}{\log\log R}\right) \tag{10.14}$$

Here R counts both jobs and bins. The asymptotic expression does not predict a finite serving queue at a given fleet size, because arrivals, departures, job sizes, and service rates are absent from this model. The power of two choices (principle 19) changes the occupancy tail using one additional probe.

10.11.3 The power of two choices

Polling every replica in a large fleet on every arrival turns the load balancer into a coordination bottleneck, yet picking one replica blindly leaves hot queues unchecked. A foundational result in load-balancing theory (Mitzenmacher 2001) shows that querying only two random servers before making a routing decision provides exponentially better load distribution than random assignment in the static balls-into-bins model. The procedure is deliberately small, but the load signal must match the serving workload. The destination must atomically admit or reject the request because probe results can become stale; a rejection requires bounded retry or overload handling. For a homogeneous vision fleet, current queue length may be enough. For an LLM fleet, active tokens, estimated remaining decode work, and KV-cache pressure better approximate the work a replica has already accepted. Algorithm 10.4 states a representative routing loop: probe two replicas, compare normalized work, and route without polling the whole fleet.

Algorithm 10.4 Power-of-two choices request routing

Require: eligible replica set $\mathcal{S}$; request r; positive capacity weights w_s; per-replica work signal W_s (queue length, active tokens, predicted work, or KV-cache occupancy); finite attempt budget
Ensure: an admitting replica for r, or an overload rejection
for each allowed attempt **do**
  if $\mathcal{S}$ is empty **then**
    return overload rejection
  end if
  sample two candidate replicas s_a, s_b from $\mathcal{S}$ with replacement ▷ *capacity-weighted*
  query W_{s_a}, W_{s_b}; estimate incremental work $\hat{c}(r)$ (expected output tokens or KV pages)
  $u_s \leftarrow (W_s + \hat{c}(r))/w_s$ for $s \in \{s_a, s_b\}$ ▷ *post-admission load*
  $s^\star \leftarrow$ candidate with smaller u_s ▷ *ties: cache affinity*
  if $s^\star$ atomically admits r and reserves its local capacity **then**
    return $s^\star$
  end if
end for
return overload rejection

Two probes avoid an all-replica scan. The practical algorithm uses predicted work and capacity weights; it is a serving heuristic, not the uniform static placement process. For the latter process, choosing the less occupied of two independently sampled bins gives equation 10.15:

$$\mathbb{E}[\text{max bin occupancy}]_{\text{two choices}} = \Theta(\log\log R) \tag{10.15}$$

Theorem 10.2: Power-of-two choices load-balancing bound

For R equal-size jobs placed sequentially into R initially empty bins, without departures, uniform random placement has maximum occupancy $\mathcal{O}(\log R/\log\log R)$ with high probability. Choosing the less occupied of two independent uniform candidates reduces this to $\mathcal{O}(\log\log R)$.

The bound shows that limited probing substantially improves maximum occupancy in this static model. Two probes avoid the R-probe cost of examining every bin. Applying the idea to serving requires a load signal that reflects remaining work and available capacity; the theorem does not establish near-optimal latency for changing queues, heterogeneous replicas, or stale measurements.

The mechanism is the suppression of repeated unlucky placement. If m of R bins have occupancy at least k, one uniform probe hits that set with probability m/R, whereas two independent probes both hit it with probability $(m/R)^2$.

10.11.4 Weighted and adaptive load balancing

Sending equal request rates to replicas with different hardware capacities saturates the slower tier while leaving the faster tier underutilized. In a mixed A100 and T4 fleet, the lower-capacity replicas reach saturation first as demand increases; equal routing does not imply overload at every arrival rate. Weighted assignment accounts for measured capacity by allocating the following long-run request share to each server:

$$\Pr(\text{route to server } i) = \frac{w_i}{\sum_j w_j}$$

where w_i is the weight, or capacity, of server i. Weighted two-choices applies the same idea to tail control: sample two servers with probability proportional to their weights, compare current load relative to capacity, and route to the lower relative load. The weighted-routing example makes this policy concrete for mixed GPU hardware.

Example 10.8: Heterogeneous GPU cluster

Scenario: A serving cluster mixes 10 H100 GPUs (1000 QPS each) and 20 A100 GPUs (600 QPS each), processing 15,000 QPS target traffic.

Diagnosis: Unweighted round-robin routes equal request counts (500 QPS) to all nodes, leaving H100s underutilized (50 percent) while leaving less latency headroom on A100s (83.3 percent). Weighted routing assigns H100s a weight of 4.5 percent and A100s a weight of 2.7 percent, balancing utilization across tiers (68.2 percent vs. 68.2 percent).

Systems lesson: Heterogeneous inference clusters require capacity-weighted routing. Equal request volumes produce unbalanced load when hardware tiers have different service rates.

Static weights handle known capacity differences, but production replicas also change over time. Adaptive load balancing adjusts weights dynamically based on observed performance:

```
For each server i:
    latency[i] = exponential_moving_average(observed_latency)
    weight[i] = 1 / latency[i]  # Inverse latency weighting
```

Inverse latency weighting reduces traffic to replicas whose observed latency rises. That observation mixes server condition with request cost: a replica receiving longer generations can appear slow even when its hardware is healthy. Delayed measurements can also make weights oscillate as traffic shifts. The heuristic therefore needs evaluation against comparable request classes, measurement lag, and overload behavior before its weights can be treated as capacity estimates.

10.11.5 Least-connections load balancing

Even capacity-weighted routing can overload a replica that draws several long-running generations in a row while neighboring servers finish short replies. Least-connections routes to the server with the fewest active requests, incrementing a counter on dispatch and decrementing it on completion. This accounts for unfinished requests that round-robin does not inspect, though a connection carrying multiple requests needs different bookkeeping. The policy improves routing only to the

extent that the count represents the work and capacity competing at each replica. Connection counts are themselves a proxy. Two in-flight LLM requests can differ fiftyfold in remaining decode work, from a 10-token reply to a 500-token essay, so active tokens or estimated remaining work track load more faithfully than the count of open requests.

The challenge is maintaining accurate connection counts in a distributed system. A centralized counter gives a single source of truth but can become a bottleneck. Distributed counters with gossip scale better but may route using stale information. Sampled least-connections combines the idea with two choices by probing only a subset of servers and choosing the minimum.

No routing policy wins on every axis. Random and round-robin assignment have low decision overhead but ignore work in progress. Two-choice routing samples load state, while least-connections tracks outstanding requests at the cost of maintaining current counts. Neither counts nor limited probes fully describe remaining work when request costs vary widely. Which policy meets a tail target is an empirical question, evaluated with the same arrival trace, request-length mix, and replica capacities.

10.11.6 Consistent hashing for stateful routing

Routing each turn of a multi-turn conversation to a different replica forces the receiving server to recompute or fetch the entire conversation's KV cache before generating a single new token. How much state a serving system retains across requests determines what must be preserved when traffic moves between replicas. Vision inference and individual embedding lookups can expose stateless interfaces because they need no session history from an earlier request. A ready replica with the required model and data can then accept a rerouted request. Adding replicas still requires loading that state, distributing traffic, and checking downstream capacity; failure recovery also requires spare capacity and a policy for retrying interrupted work. Statelessness removes session-state migration from these operations, but does not remove their resource costs.

↩ Stateless and stateful failure recovery mechanics are established in section 7.11.1.

An autoregressive request is stateful while it is generating because its KV cache belongs to that in-flight sequence. Cross-request statefulness is optional. An API may send the full conversation history with every turn and rebuild KV state during prefill, making any replica eligible for the next request. A serving platform may instead retain a completed turn's KV state as a session cache, reducing repeated prefill at the cost of replica affinity, cache lifetime management, and recovery complexity. Sticky routing is an optimization for that retained cache, not a correctness requirement of autoregressive models.

The scaling divide is therefore an interface and cache-policy decision. Every live generation needs request-scoped state, but the service decides whether state survives the request boundary. Session-cached service saves prefill work but must drain, migrate, or rebuild cached state during scaling and failure, so affinity earns its place only when measured cache-reuse value exceeds the resulting imbalance and recovery costs. Recommendation sessions that carry user context and streaming inference that depends on state from earlier frames benefit from the same routing affinity, while a stateless LLM API may still use prefix-aware routing when a reusable prefix happens to be cached.

Consistent hashing[6] (Karger et al. 1997) maps requests to servers based on a hash of the routing key (user ID, session ID). Servers and keys are hashed onto the same circular space, called the ring, and each key belongs to the first server reached by moving clockwise from the key's position:

$$\text{server}(\text{request}) = \arg\min_{s \in \mathcal{S}_{\text{srv}}} \text{distance}(\text{hash}(\text{key}), \text{hash}(s))$$

where $\mathcal{S}_{\text{srv}}$ is the set of servers mapped onto the ring and distance is measured clockwise from the key's hash.

For fixed ring membership, consistent hashing sends the same key to the same server. Membership changes affect only part of the ring, because a server that joins or leaves changes the owner of only the arc of keys next to it. With balanced equal-capacity assignments and uniformly distributed keys, adding a server moves an expected $1/(N+1)$ fraction of keys, those in the arc it takes over, while removing one affects its expected $1/N$ share. Virtual nodes help smooth uneven ranges, but popular keys can still concentrate traffic even when key counts are balanced. For session-cached LLM serving, affinity increases the chance of reaching reusable state. That benefit must outweigh additional waiting; overload routing, failure, eviction, and remapping can still require repeated

[6] **Consistent hashing**: The ring-based algorithm was introduced for distributed caching (Karger et al. 1997) and then used in large-scale content delivery. Its limited-remapping property made it useful wherever membership changes should preserve cache locality.

prefill. The same stateful routing machinery becomes a correctness boundary when cached state is user-specific.

War Story 10.1: When cache state crossed users (2023)

Context: ChatGPT used Redis Cluster through the redis-py client to cache user information and conversation titles in a high-volume serving system (OpenAI 2023).

Mechanism: Under an Asyncio connection-pooling race condition, canceled requests left unread responses in recycled connection buffers, causing subsequent requests to receive cache data belonging to previous users.

Impact: On March 20, 2023, OpenAI took ChatGPT offline after a bug allowed some users to see another active user's conversation titles. Further investigation found possible exposure of payment-related information for a subset of active Plus subscribers during a nine-hour window; full credit card numbers were not exposed.

Fix: OpenAI sent the Redis maintainers a patch for the redis-py cancellation path and added redundant checks that the data returned from the cache belongs to the requesting user.

Systems lesson: Stateful inference infrastructure needs isolation guarantees around caches, connection pools, and routing keys. A cache is part of the serving correctness boundary, not merely an optimization.

10.11.7 Request routing for sharded models

↪ Shared prefix caches as a cross-tenant leakage channel, and the tenant-scoped caches that close it, are developed in section 13.6.3.

Treating individual GPUs inside a sharded model as independent routing targets breaks execution: a tensor-parallel group must run every layer in lockstep across all its ranks, and a pipeline group can admit new requests only at its first stage. A sharded replica is therefore one routing destination even though several devices execute the request. A tensor-parallel group presents one destination whose latency already includes its per-layer collectives (figure 10.16). A pipeline group admits requests at its first stage (figure 10.17), so the router must account for the group's service capacity rather than treat each stage as a replica. In expert-parallel architectures, the gating network (section 10.10.4.2) routes tokens across internal expert destinations after the external router has chosen the replica, so expert imbalance and replica imbalance require different controls. The result is a two-tier topology (figure 10.18) in which the external load balancer sees each shard group as one destination and the group's NVLink collectives stay inside it.

10.11.8 KV cache migration between replicas

When a replica fails or consistent hashing remaps a session to a new shard group, the new group cannot generate the next token until it either receives the conversation's gigabyte-scale KV cache over the network or reruns prefill across the entire history. Routing a request to a shard group fixes where its KV state lives (section 10.10.2), so moving an active session between groups forces a choice between network transfer and recomputation (section 10.11.6). Transfer requires a reachable source or a saved copy; after an unrecoverable replica failure, the service rebuilds from authoritative conversation history:

```
Migration options:
1. Rebuild: Re-run prefill on new replica (cost grows with context length)
2. Transfer: Send KV cache over network (cache bytes / link bandwidth)
3. Hybrid: Apply the threshold below when the source is reachable;
   otherwise rebuild

Decision threshold:
if cache_size_bytes/network_bandwidth < prefill_time:
    transfer()
else:
    rebuild()
```

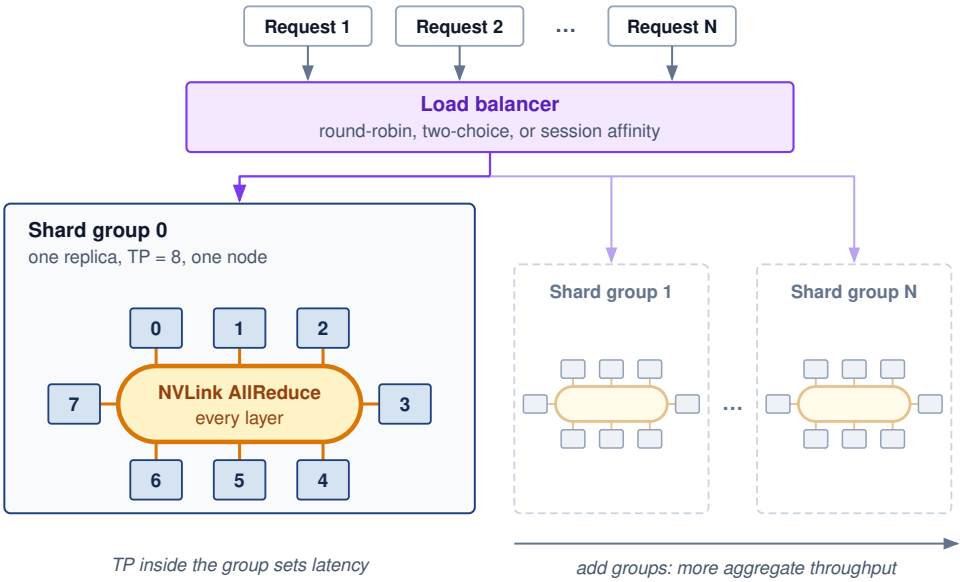

Figure 10.18: Horizontal Scaling with Shard Groups: The load balancer treats each shard group as a single logical inference server and chooses among groups by round-robin, two-choice, or session-affinity routing. Within a group, eight GPUs hold the tensor-parallel weight shards and join partial results with an NVLink AllReduce inside every layer, which the router never sees. Adding ready groups raises aggregate capacity; the group's internal parallelism affects per-request latency.

For Llama 3 70B with GQA and a 4,096-token context, KV cache is about 1.3 GB total per sequence, or about 168 MB per device under 8-way tensor parallelism. At 100 Gb/s (12.5 GB/s), the bandwidth-only estimate for one shard is roughly 13 ms, while sending the full cache over one such link takes about 107 ms. The prefill approximation counts 5.8×10^{14} FLOPs ($2 \times 7.1 \times 10^{10}$ parameters $\times$ 4,096 tokens), giving roughly 180 ms on the tensor-parallel group at the assumed 40 percent utilization. Transfer is faster in this comparison, but achieving the per-shard time for the whole group requires concurrent links with sufficient aggregate bandwidth. Shared network bottlenecks, source availability, cache conversion, and destination queueing can change the decision. Long-context attention adds rebuild work omitted from this approximation, so measured transfer and prefill times should determine the policy.

The shard group boundary also defines readiness. A group can accept work only when the devices, model state, and cache capacity required for its execution plan are available, so the load balancer needs a health signal for the serving replica as a whole.

10.11.9 Health checking and failover

An inference worker whose HTTP server responds `200 OK` to a shallow ping can still time out every real query if its weights are still loading, its KV pool is exhausted, or its CUDA context is wedged. Load balancers can route around failure only if the health signal matches the failure mode across three probe tiers (liveness, readiness, and inference correctness) that escalate from cheapest to most realistic. On an accelerator replica, each tier must test a different part of the serving state: readiness includes admission capacity as well as loaded weights and completed warmup, while an inference probe compares a fixed prompt's output or logits against expected bounds. Passing that probe does not establish correctness or capacity for every request shape, so production errors and latency remain necessary health signals. Failover then depends on the detected failure, available replacement capacity, and whether affected KV state can be transferred or must be reconstructed.

↩ Liveness, readiness, and inference-correctness probes, the failover time their settings bound, and stateful failover are established in section 7.11.3.

10.11.10 Circuit breakers and backpressure

When a GPU inference server slows down, routing more requests to it can turn a local problem into a fleet-wide failure. Consider an illustrative rise in service time from 50 ms to 500 ms. Queues

accumulate, client timeouts generate retries, and rerouted work consumes neighboring replicas' headroom. The cause might be a stalled dependency, device degradation, or resource contention. The load balancer needs a policy for withdrawing an unhealthy replica while controlling the displaced load.

The circuit breaker pattern (Nygard 2007) limits repeated work sent to an unhealthy destination. In the closed state, the server receives normal traffic. When the configured error, timeout, or latency threshold is exceeded, the breaker opens and requests fail fast or become candidates for rerouting. After a recovery interval, a half-open state admits limited probes, closing on recovery or reopening on failure (figure 10.19). Opening a breaker does not create replacement capacity, so rerouting must respect admission limits and retry budgets or it can overload the surviving replicas.

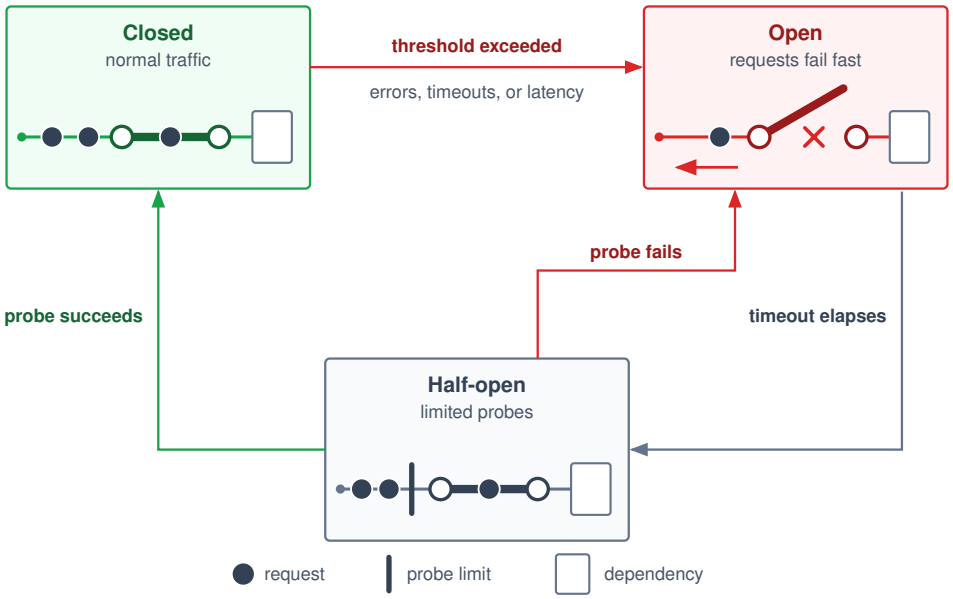

Figure 10.19: Circuit Breaker States: The circuit breaker protects the system from cascading failure. Closed: normal operation. Open: error threshold exceeded; all requests fail fast to prevent resource exhaustion. Half-open: after a timeout, a limited number of requests are allowed through to probe the dependency's health. Success resets to closed; failure returns to open.

Backpressure propagation complements the breaker by making overload visible upstream. When queue depth crosses a threshold, the server returns a signal such as `503 Service Unavailable`; the load balancer marks the replica as degraded and routes fewer requests to it. If all replicas become degraded, the system shifts from routing control to admission control, rejecting some requests at the edge rather than accepting work it cannot serve within the latency SLO. Load-shedding policies decide which requests to reject first. Circuit breakers isolate unhealthy replicas; backpressure tells the rest of the system to slow down before retries amplify the failure.

Backpressure works only when the request contract propagates through the whole serving graph. The gateway converts the client timeout into an absolute deadline and every downstream stage subtracts elapsed time before admitting more work. If the client disconnects or the deadline expires, cancellation must reach queued CPU work, accelerator batches, retrieval calls, and streamed output. A canceled request that remains in a continuous batch is zombie work: it consumes KV pages and decode bandwidth while producing no useful response. Some kernels cannot be interrupted mid-launch, so cancellation commonly takes effect at the next batch or token boundary; the scheduler must measure that cancellation lag rather than promise instantaneous preemption.

↩ Load shedding, model fallback, and the other strategies that trade quality for availability are established in section 7.12.2.

The same request identifier must connect routing, queueing, prefill, decode, cache allocation, and response streaming. A useful serving trace records admitted and rejected requests, queue delay, TTFT, TPOT, input and output tokens, cache pages, batch membership, preemptions, cancellations,

and fallback decisions. Aggregate accelerator utilization cannot distinguish useful work from late or canceled work. SLO goodput and the wasted-token rate close that gap. SLO goodput reports how much completed work met its contract. The **wasted-token rate** is the fraction of generated tokens spent on requests that were canceled, timed out, or delivered after their deadline, the execution the system performed after the result had lost value.

A saturated LLM fleet operating under severe KV cache pressure has three workload-specific relief options. Engines can apply context eviction through an explicitly supported sliding-window or approximate policy with a quality check, or truncate retained history and recompute its state. Silently dropping old KV entries is not equivalent to recomputing a shorter conversation, because later cached states already depend on the earlier context. Disabling speculative execution reduces draft work, but reclaiming its memory also requires unloading draft weights and releasing the associated cache state. Finally, fallback routing can send eligible requests to a smaller model with available capacity, provided that model meets the application's fallback quality contract.

These options change latency, resource use, or model behavior in different ways. The application must specify which changes are acceptable at each pressure tier. Context truncation can remove information essential to an answer, and a smaller model can fail the task's quality requirement. In those cases, an explicit rejection can be preferable to an apparently successful but unacceptable response.

Routing relies on estimates of replica capacity, but shared hardware can make that capacity change when another workload becomes active. Critical production models may share a cluster or device with experimental models and developer endpoints. Protecting the serving contract under that interference requires multi-tenant isolation in addition to routing.

10.12 Multi-Tenancy and Isolation

Consider a critical customer-support chatbot running on the same GPU server as an experimental computer vision script. If the experimental workload monopolizes the PCIe bus or memory bandwidth, the chatbot can miss its latency target. Sharing saves capacity only when the resulting interference remains within each service's contract. Fleet quotas, fair-share hierarchies, and priority preemption decide how much of the cluster a serving system receives; the problem here is controlling contention among the tenants who share that allocation, request by request.

↩ Fleet quotas, hierarchical fair-share, and the priority classes that govern preemption are established in section 8.11.

Sharing a serving pool can raise utilization and reduce separately provisioned capacity when tenant demand peaks occur at different times. This statistical-multiplexing gain weakens when bursts are correlated, and shared resources introduce interference even when aggregate demand fits. Pooling therefore needs both a demand model and enforceable isolation.

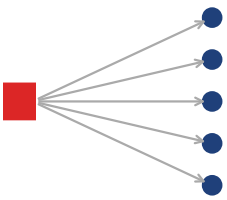

One tenant's burst can degrade every workload sharing the pool.

A burst from one tenant can delay others through shared KV-cache capacity, network bandwidth, and accelerator time. The pool must preserve each tenant's data boundary and service contract; different SLOs can require different admission limits, reservations, or placements. Noisy-neighbor interference comes first because it connects shared-resource contention directly to those latency contracts.

10.12.1 Noisy neighbor problems

A noisy neighbor, the tenant whose consumption of a shared resource slows the others on the same hardware, interferes inside a serving pool through three resources at once, and each is denominated in units a request-level scheduler can meter. Unchecked demand for KV-cache capacity, network egress bandwidth, or accelerator execution time lets one bursty workload erode the headroom of every co-located tenant.

GPU memory contention can limit admission before compute capacity is exhausted. Consider three tenants sharing a 60 GB KV pool with nominal 20 GB shares. Without enforced limits, one tenant's long-context requests can raise its allocation to 45 GB. In the scenario, the others lose capacity from 200 concurrent sequences each to 75, a 62.5 percent reduction. That is a concurrency loss, not a proportional throughput prediction; the effect depends on their execution regime and preemption costs. Egress bandwidth and accelerator time are additional shared resources whose contention must be measured alongside memory occupancy.

↩ Noisy-neighbor contention and the device-level interference channels of shared accelerators (memory bandwidth, cache hierarchy, host-device interconnect, and thermal dissipation) are established in section 8.10.1.

Example 10.9: Quantifying noisy neighbor impact

Scenario: An inference platform serves 10 tenants on shared H100 GPUs at 100 QPS each under a 50 ms p99 SLO. Tenant 3 bursts to 500 QPS, 5× its baseline.

Diagnosis: Without quotas, contention raises pool-wide p99 into the 22 ms to 45 ms range in this scenario, eroding every tenant's headroom while remaining within the stated target. With quotas, Tenant 3 is limited to 120 QPS and the other tenants retain 15 ms p99. Throttling turns Tenant 3's excess into rejected or delayed requests, which are judged against Tenant 3's own contract rather than spread across the pool.

Systems lesson: Shared capacity raises utilization only when it is coupled with isolation. Quotas turn pool-wide erosion of latency headroom into an admission decision confined to the tenant that caused it; without them, a larger burst or a tighter SLO turns the same erosion into violations across the pool.

10.12.2 Resource quotas and fair sharing

Limiting tenants by request count alone fails when one request carries a 20-token query and the next submits a 100,000-token document that monopolizes KV-cache pages and prefill compute. Resource quotas limit what each tenant can consume across every scarce resource, preventing any single tenant from monopolizing shared hardware. Hard quotas turn shared serving into admission control: request rate and concurrency bound arrival pressure, while input tokens, output tokens, KV-cache pages, accelerator time, and egress bytes bound the work performed after admission. Listing 10.1 shows representative gates so a tenant can be throttled at the scarce resource boundary instead of degrading the whole batch.

Listing 10.1: Multidimensional Resource Quotas: A snapshot check compares prospective resource use with tenant limits. Atomic reservation and scheduler enforcement are separate requirements.

```
class TenantQuota:
    max_concurrent_requests: int  # e.g., 100
    max_kv_cache_mb: int  # e.g., 20,000
    max_qps: int  # e.g., 1,000
    max_input_tokens_per_minute: int
    max_output_tokens_per_minute: int
    max_batch_tokens: int  # e.g., 50,000


def check_quota_snapshot(tenant_id, request):
    quota = get_quota(tenant_id)
    usage = get_usage(tenant_id)

    if usage.concurrent + 1 > quota.max_concurrent_requests:
        return RateLimitError("concurrent request limit")
    if usage.kv_cache_mb + request.reserved_kv_mb > quota.max_kv_cache_mb:
        return RateLimitError("memory limit")
    if usage.requests_last_second + 1 > quota.max_qps:
        return RateLimitError("rate limit")
    if usage.input_tokens + request.input_tokens > quota.max_input_tokens_per_minute:
        return RateLimitError("input-token limit")
    if usage.output_tokens + request.max_output_tokens > quota.max_output_tokens_per_minute:
        return RateLimitError("output-token limit")

    if request.prefill_chunk_tokens > quota.max_batch_tokens:
        return RateLimitError("prefill chunk exceeds batch-token limit")

    return None  # Snapshot is eligible; no resources have been reserved.
```

This listing assumes that `reserved_kv_mb` includes the requested prompt and output-token budget, and that usage counters include outstanding reservations in their stated time windows. A

passing snapshot is not an admission guarantee: simultaneous callers can observe the same free capacity. Production admission must check and reserve resources atomically, roll back failed admissions, and release reservations on completion or cancellation. The batch scheduler separately enforces the aggregate batch-token limit. Quotas protect the resources they account for; shared compute, networking, and control-plane contention still require isolation and measurement.

Soft quotas permit borrowing unused capacity while preserving an enforceable share when demand rises. The unit of account matters: equal request counts can hide large differences in prefill work, output length, and retained state. Input and output tokens therefore need separate accounting, with memory and concurrency limits alongside them. The progressive-filling approach in section 8.4.1 reduces to max-min fairness when only one divisible resource constrains the allocation. For three tenants demanding 30,000, 20,000, and 80,000 output tokens/sec from an idealized 100,000-token/sec pool with equal token costs, the allocation is 30,000, 20,000, and 50,000. Coupled prefill, decode, and memory constraints require joint feasibility checks; independent fair allocations do not automatically compose into a feasible service.

10.12.3 Priority scheduling

First-come, first-served scheduling treats a paid interactive query and a background batch job identically, letting a surge of free-tier or offline traffic delay revenue-critical requests past their SLO. Priority-aware scheduling gives eligible high-priority requests precedence at scheduling boundaries: a production API might classify revenue-generating customer requests as critical, internal batch processing as standard, and free-tier traffic as best-effort. Already-running kernels and shared resources can still delay critical work, so admission limits and reserved capacity must complement priority ordering. Priority becomes an isolation tool once its classes map to tenants and service tiers. Fleet priority classes govern preemption across allocations; request priority applies inside a serving allocation. Table 10.19 illustrates a policy that combines precedence with reservations. Those reservations must be provisioned and enforced separately: priority alone cannot create capacity, interrupt every running kernel, or guarantee a deadline under unbounded critical traffic.

Table 10.19: Priority Scheduling Classes: Three-tier priority hierarchy for multi-tenant inference. Critical traffic gets 100 percent reserved capacity and can preempt any lower class; standard traffic shares capacity proportionally and can preempt best-effort; best-effort runs on leftover capacity with no guarantees.

Class	Use case	Preemption	Resource guarantee
Critical	Revenue-generating	Can preempt lower	100% reserved
Standard	General traffic	Can preempt best-effort	Weighted share
Best-effort	Background, batch	Cannot preempt	No guarantee

Preemption extends priority to requests already in flight, but priority does not make state movement free. The retain, swap, recompute, or reject choice of section 10.7.7 still decides how a displaced request's KV state is handled; priority adds an input to that choice rather than replacing it. A priority class determines who may be displaced, while remaining service time, state size, and deadline slack determine which victim is cheapest. Ignoring those costs lets aggressive preemption consume more capacity than it releases.

10.12.4 Bulkhead pattern

Quotas and priority queues inside a single shared serving process cannot stop a malformed request or out-of-memory fault from crashing the worker and dropping every tenant's in-flight generations at once. The bulkhead pattern[7] (Nygard 2007), named after ship compartments intended to contain flooding, partitions resources across separate concurrency pools, processes, or dedicated hardware so exhaustion or failure in one allocation stays contained.

Dedicated replicas separate more resources than request limits inside one process. Figure 10.20 assigns distinct compute and memory pools to tenant groups, so an engine crash or exhausted local allocation in the standard tier need not stop the gold-tier engines. Shared networking, storage, and control-plane services can still carry interference across that boundary.

[7] **Bulkhead limits**: A divider isolates failure only if it closes the relevant boundary. The Titanic's incomplete bulkheads allowed water to cascade between compartments; similarly, a soft quota cannot isolate shared memory or bandwidth that remains globally exhaustible.

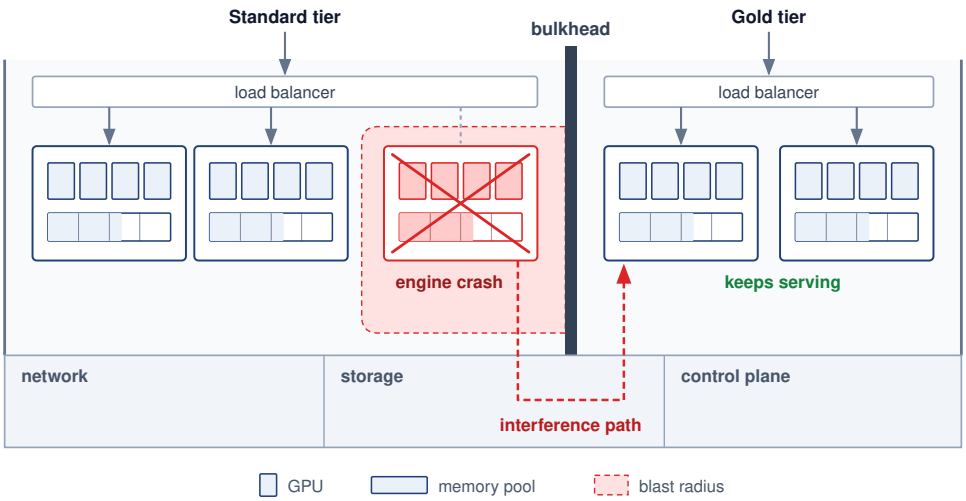

Figure 10.20: Bulkhead Isolation Patterns: To prevent cascading failures in multi-tenant systems, bulkheads isolate resources. Deployment-level bulkheads (shown) assign dedicated physical replicas to high-priority tenants, separating their compute and memory pools. Request-level bulkheads enforce strict concurrency limits within shared processes. Shared networking, storage, and control-plane dependencies remain possible paths for interference.

Deployment-level bulkheads cost pooling efficiency and the operational overhead of dedicated infrastructure. Request-level bulkheads are cheaper but weaker. Caps such as 8,000 input tokens, 2,000 output tokens, and 30 seconds of execution bound request size and duration when the scheduler enforces them, but they still need concurrency limits, reservations, and cancellation handling, and they cannot contain a crash. A malformed request that crashes a shared worker interrupts every tenant on that worker, so service tiers that require stronger fault isolation need placement boundaries to match.

Example 10.10: Bulkhead configuration for API tiers

Scenario: An LLM API platform serves Enterprise, Professional, and Free service tiers across shared and dedicated hardware.

Diagnosis: Unpartitioned pools allow Free-tier bursts or shared-worker failures to affect Enterprise requests. The example assigns Enterprise to dedicated GPU pools with a 99.9 percent availability target, allocates 70 percent of the remaining shared capacity to Professional, and caps Free at 30 percent best-effort capacity. Meeting that target still requires reliability evidence for the dedicated pools themselves.

Systems lesson: Bulkheads concentrate isolation overhead where business requirements dictate. Hardware partitioning reduces cross-tenant compute and memory contention. SLO compliance still depends on sufficient capacity, admission policy, shared dependencies, and failure handling.

Multi-tenant isolation limits the blast radius of excess demand, but it does not create capacity. Once protected queues approach their latency budgets, the service needs a control loop that can add replicas before admission begins to fail.

10.13 Autoscaling and Capacity Control

A serving autoscaler that reacts only to average accelerator utilization can scale in the wrong direction. Decode can saturate memory bandwidth while arithmetic units appear idle, long prompts can consume prefill capacity before request count rises, and queued work can violate an SLO while

fleet-wide utilization remains moderate. The control signal must represent outstanding work relative to phase-specific service capacity. Every replica the controller adds is also a claim on accelerators that training jobs may hold, so fleet-level arbitration with training bounds how quickly and how far serving can grow.

↩ The arbitration between serving replicas and training jobs for shared accelerators, including the minutes-long cost of switching a device between them, is established in section 8.10.3.

For a conventional fixed-cost model, queueing delay and request rate per ready replica can guide scaling. Autoregressive serving also needs prompt-token and output-token rates, KV occupancy, and measured phase service rates. A colocated replica must satisfy all phase requirements; disaggregated prefill and decode pools need separate counts and a feasible transfer budget. SLO goodput measures useful supply, but rejected work and backlog must remain visible as unmet demand. A latency threshold can detect deterioration after a model change, yet controller timing and headroom still need validation. Adding replicas also depends on their startup delay, and resizing a running replica's accelerator allocation may require a restart or state migration, depending on the runtime and execution plan.

Scale-up and scale-down require asymmetric evidence. Scale-up acts on sustained queue growth, forecast demand, or declining deadline slack because cold-start delay makes late reactions expensive. Scale-down requires low demand over a longer window, enough spare capacity to survive one replica failure, and a drain plan for in-flight requests and any optional session cache. This hysteresis prevents oscillation in which replicas repeatedly load large weights, accept work, and terminate before amortizing startup cost.

The controller also needs observability that distinguishes demand from impaired supply: the per-request serving trace of section 10.11.10, plus ready replicas, model-load failures, and per-phase service rate. Without readiness and service-rate measurements, an autoscaler can respond to a sick replica by adding cold replicas indefinitely while the original failure remains hidden.

10.13.1 Scaling response time

If a traffic spike hits a serving pool in 10 seconds but acquiring a GPU node and loading 140 GB of weights takes five minutes, the autoscaler cannot prevent the queue from overflowing during the gap. The scaling response budget is the sum of the control loop's detection, decision, provisioning, and warmup delays, as equation 10.16 makes explicit:

$$T_{\text{response}} = T_{\text{detect}} + T_{\text{decide}} + T_{\text{provision}} + T_{\text{warmup}} \tag{10.16}$$

Each term has its own typical duration and its own primary lever, from a shorter metrics interval for detection to warm pools for provisioning and precompilation for warmup (table 10.20).

Table 10.20: Scaling Response Time Components: The four phases of a scaling event with illustrative durations and their primary optimization lever. Provisioning dominates this example; a deployment must measure its own critical path, which may instead be dominated by artifact loading.

Component	Duration	Optimization
Detection	10–60s	Reduce metrics interval
Decision	1–5s	Faster autoscaler
Provisioning	30s–5min	Warm pools
Warmup	5–30s	Precompilation

Provisioning can dominate the illustrated budget, making already allocated capacity valuable. The large-model example that follows shows that artifact loading can instead dominate startup; the relevant term is the one measured on the deployment's critical path. Finer-grained metrics can reduce detection delay but increase noise, while prepared scaling decisions and precompiled engines reduce other costs. A warm pool removes only the phases already completed at its stated readiness level.

10.13.2 Cold start and predictive scaling

When a 70B replica takes nearly seven minutes to boot and stream its weights from object storage, waiting for a morning traffic ramp to appear in queue metrics guarantees several minutes of over-

loaded replicas. A cold start for a 70B replica shows the scale of that delay, and a demand forecast shows how a controller can begin the wait before the traffic arrives.

Example 10.11: Cold start timeline for Llama-70B

Scenario: An autoscaling engine provisions a Llama-70B replica during a traffic surge, using an H100-based allocation large enough for its FP16 weights, cache, and runtime state.

Diagnosis: The scenario's sequential startup estimate is 6.7 minutes ($T_{\text{cold-start}} = T_{\text{provision}} + T_{\text{load}} + T_{\text{warmup}}$): 1 minute for provisioning, 4.7 minutes to fetch 140 GB of artifacts at the assumed 0.5 GB/s, and 1 minute for device loading and warmup. Cached artifacts or overlapping phases can shorten this path; scarce capacity and shared storage can lengthen it. If existing headroom cannot absorb the surge, newly requested replicas arrive too late to protect short deadlines.

Systems lesson: Cold-start latency depends on artifact size, effective storage bandwidth, provisioning delay, compilation, and warmup. The controller must measure this distribution and begin provisioning far enough ahead that its high percentile fits inside the forecast horizon.

Example 10.12: Predictive scaling for traffic

Scenario: A commercial chatbot service faces predictable diurnal traffic patterns, ramping from 500 QPS at night to 3,000 QPS at peak.

Diagnosis: A controller that waits for the morning ramps at 06:00 and 09:00 UTC before requesting capacity must cover the assumed 7-minute startup interval with existing headroom or admission limits. Without enough ready capacity, queues grow while replicas start. Forecasting can shift that wait ahead of the expected demand.

Systems lesson: Predictive autoscaling converts demand forecasts into planned pre-warming windows (table 10.21). Starting 30 minutes ahead of the expected ramps exceeds the modeled startup time. Readiness still depends on actual provisioning and warmup, and forecast errors can leave the fleet short or idle. Reactive correction and an overload policy remain necessary.

Table 10.21: Predictive Scaling Schedule: Per-time-bucket scaling actions, active replicas, and replicas in cold-start warming for the chatbot workload.

Time	Action	Replicas active	Replicas starting
05:30	Scale up	5	+10 warming
06:00	Traffic ramp	15	-
08:30	Scale up	15	+15 warming
09:00	Peak traffic	30	-
17:00	Scale down	20	-10 terminating
20:00	Scale down	10	-10 terminating
00:00	Scale down	5	-5 terminating

Systems insight: Against an assumed reactive peak of 45 replicas, the forecast-based peak of 30 is a 33 percent reduction in peak replica count at equal unit cost. The reactive controller over-provisions because it reacts to a queue that has already formed and must add capacity for the backlog as well as the demand. Peak count is only part of the bill. Total spending counts replica-hours, including the pre-warming window, so the forecast pays only when the replica-hours spent warming ahead of demand are fewer than those the reactive overshoot would consume.

Example 10.12 shows how capacity can be prepared before demand arrives. Figure 10.21 depicts a forecast that anticipates the surge, while the reactive curve starts later and temporarily lacks capacity. Whether that deficit produces queueing, rejection, or missed deadlines depends on admission and existing headroom. The predictive curve illustrates a successful forecast, not a guarantee that future demand or startup time will match it.

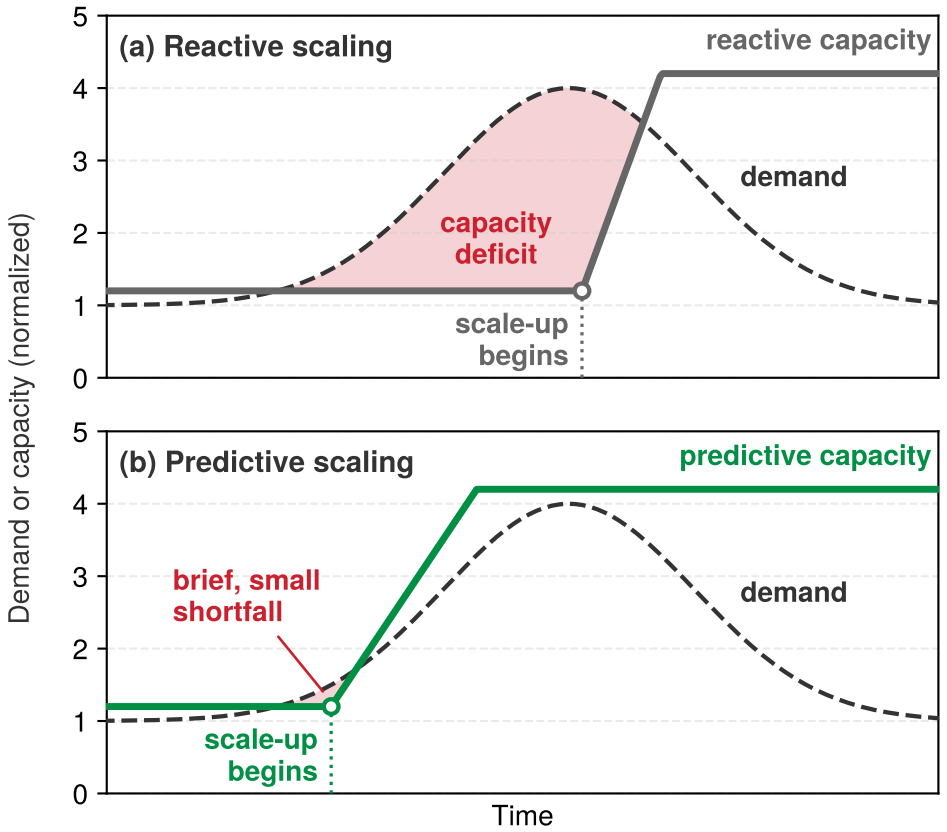

Figure 10.21: Predictive vs. Reactive Scaling: Reactive scaling (gray line) responds to traffic spikes after they occur, leaving a capacity deficit (red region) during cold start in this scenario. Predictive scaling (green line) anticipates traffic and begins provisioning before the spike peaks, shrinking that deficit to a brief, small shortfall at onset in the illustrated forecast. Its early ramp does not eliminate every interval when demand exceeds capacity.

10.13.3 Warm pool management

Demand forecasts still miss unscheduled traffic spikes, and cold-starting a large model from scratch takes minutes that an interactive SLO cannot spare. Maintaining prewarmed replicas in a warm pool buys readiness before unexpected demand arrives. Sizing that pool begins with the total capacity needed during the provisioning delay, then subtracts capacity already ready to serve. If peak demand doubles and the chosen headroom factor is 1.5, the target is three times baseline capacity, assuming unchanged per-replica throughput. Partly prepared replicas count only if they can become ready before that capacity is needed. The expression Pool size $\times$ GPU cost/hour $\times$ Idle fraction estimates the hourly cost attributed to idle GPU capacity; it excludes busy capacity and non-GPU resources. Readiness comes in tiers, each trading activation delay against the cost of waiting (table 10.22).

Table 10.22: Tiered Warm Pool Readiness Levels: Readiness trades activation delay against idle resource cost. Actual delay and billing depend on the deployment. Cold capacity is subject to quotas and physical availability.

Tier	State	Activation requirement	Cost while waiting
Hot	Model running on allocated GPU	Route admitted work	Allocated GPU and service resources
Warm	Capacity allocated; runtime partly prepared	Finish initialization and validation	Allocated resources remain billable
Cold	GPU not allocated	Acquire capacity, load, and warm up	Artifact storage; no allocated GPU

When a traffic spike hits, the controller routes work to ready replicas, activates prepared capacity, and requests additional instances if needed. Pool sizing must cover the measured provisioning delay and include an overload policy for capacity that cannot be acquired in time.

10.13.4 Spot and preemptible instances

Revoking a spot GPU replica mid-generation destroys every in-flight sequence's KV cache unless the worker can drain or migrate its active requests inside the cloud provider's two-minute warning window. Discounted interruptible capacity enters serving through the same purchasing tiers that fleet capacity planning uses (table 8.4): reserved capacity for the steady baseline, on-demand capacity for bursts, and spot capacity for work that tolerates revocation. Equation 8.3 derives the effective cost of preemptible capacity from its productive fraction. Serving adds a loss that training checkpoints do not cover, because the productive fraction must also pay for draining, rerouting, and rebuilding in-flight KV state without violating the affected requests' SLO. Listing 10.2 orders the shutdown path so new work stops first, in-flight work drains while time remains, recoverable state is persisted, and the load balancer stops sending traffic before the instance disappears.

↩ Capacity purchasing tiers and the effective cost of interruptible capacity are established in section 8.8.

Listing 10.2: Spot Termination Handling: Graceful shutdown sequence that stops admission, deregisters from the load balancer, drains in-flight requests, and persists recoverable state.

```
async def handle_spot_termination():
    # Received 2-minute warning
    # 1. Stop admission and remove the replica from routing
    stop_accepting_requests()
    deregister_from_loadbalancer()

    # 2. Complete in-flight requests (if possible)
    await complete_inflight(timeout=90)

    # 3. Save state for resumption elsewhere
    save_kv_cache_to_storage()

    # 4. Terminate gracefully
    shutdown()
```

Replica-level draining protects in-flight requests only if the upstream router avoids placing strict-SLO traffic on interruptible capacity in the first place. The gateway in figure 10.22 separates traffic by its tolerance for capacity revocation, and an eviction notice triggers draining when time permits; requests that cannot finish need a permitted retry, reconstruction, or rejection path.

Savings under this split depend on provisioned GPU-hours, discounts, retries, and reserve capacity rather than on the traffic share, and spot capacity changes none of the fault-tolerance requirements of the SLO-critical path. Even a well-scaled single-region fleet, however, remains bounded by wide-area network latency to distant users and by the blast radius of a regional outage.

10.14 Global Inference Infrastructure

A fleet confined to one region can batch, shard, route, isolate, and autoscale well, yet a user in Sydney still pays roughly a 200 ms round trip to a US-East deployment before any of that machinery

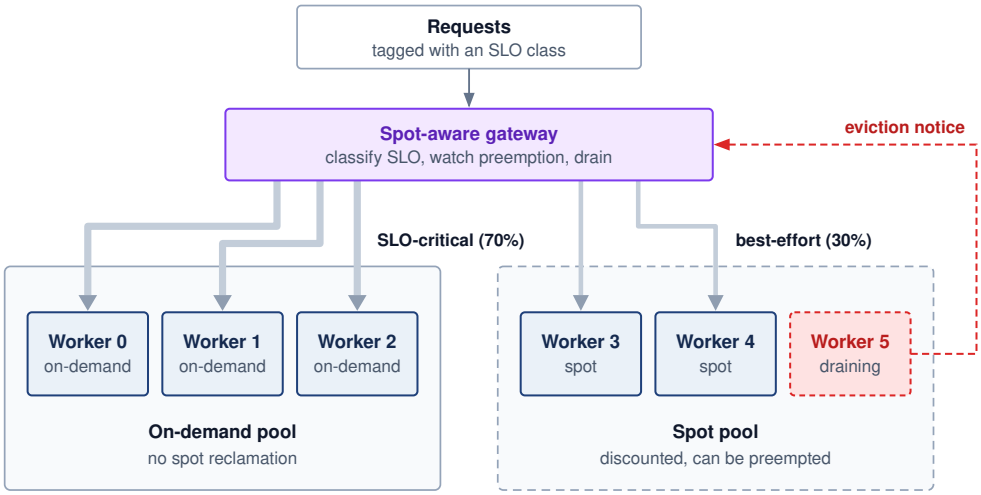

Figure 10.22: Spot-Aware Traffic Distribution: The gateway classifies requests by SLO and routes 70 percent of traffic, the SLO-critical share, to on-demand replicas that face no spot reclamation, and 30 percent, the best-effort share, to discounted spot replicas that can be preempted. When a spot worker receives an eviction notice, the gateway stops routing to it while it drains, and best-effort requests tolerate occasional retries. The traffic split is illustrative; cost savings depend on provisioned capacity, discounts, retries, and idle headroom.

acts. A regional outage also removes every replica at once. Extending resource isolation and failure boundaries beyond one data center introduces placement decisions at geographic scale. Distance, regional failures, and data-placement requirements can motivate regional replicas, caches, or a centralized service with an explicit network budget. Figure 10.23 illustrates three tiers: edge points of presence for reusable responses, regional gateways for authentication and routing, and GPU clusters for model execution. Its 10 ms, 50 ms, and 2 s budgets are scenario targets, not measured performance or universal deadlines. The application must define whether its execution budget governs first response, individual tokens, or completion of the entire request.

10.14.1 Drivers for multi-region deployment

Serving all global traffic from a single US-East cluster forces European, Asian, and Australian users to pay a 75 to 200 ms wide-area network round trip on every request and leaves the entire service vulnerable to a single regional outage. Network round-trip time (RTT) adds a latency floor for distant users, and illustrative regional timings show that floor rising with distance from the serving region (table 10.23). Actual paths, access networks, and congestion require measurement.

Table 10.23: Network RTT by User Location: A single US-East deployment imposes a 75 to 200 ms RTT on European, Asian, and Australian users, an order of magnitude above the local-region floor. Sequential calls that cross this boundary can add repeated wide-area delay. Regional replicas can reduce that exposed delay when placement policy, workload volume, and available capacity justify them.

User location	RTT to US-East	RTT to local region
New York	10 ms	10 ms
London	75 ms	10 ms
Tokyo	150 ms	10 ms
Sydney	200 ms	10 ms

Network delays compound when a request makes sequential calls across the distant boundary; multiple model calls inside one backend region do not each pay another client-to-region RTT. Regional replication can shorten that path and separate failure domains. Data-placement policies

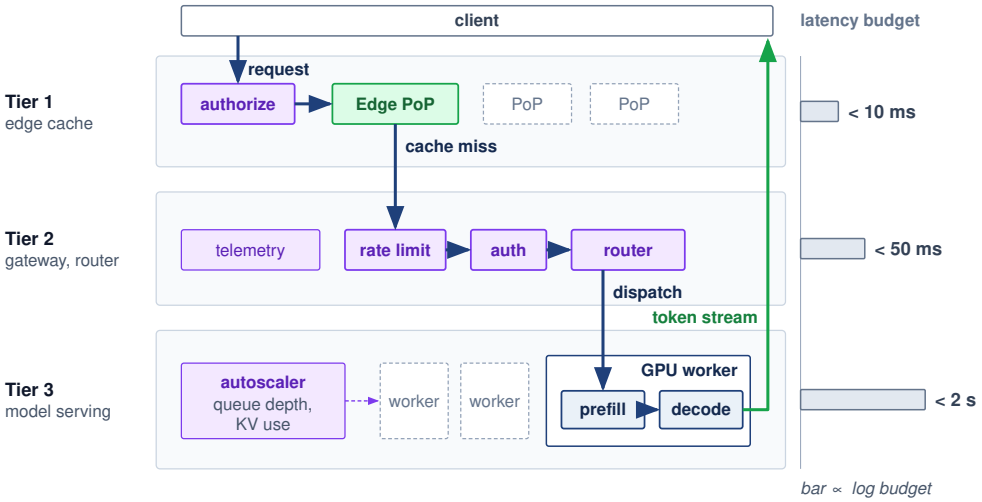

Figure 10.23: Serving Deployment Stack: A cache miss crosses three tiers before reaching a model replica. Tier 1 edge points of presence (PoPs) enforce authorization and reuse eligibility before answering repeated requests from cache; a miss passes to Tier 2, where the gateway applies rate limits and authentication and the router picks a destination; Tier 3 GPU workers run prefill and decode while an autoscaler sizes the fleet from queue depth and KV cache use. Tokens stream back up through the tiers to the client. The per-tier latency budgets on the right, drawn on a log scale, are illustrative.

also constrain eligible destinations. The General Data Protection Regulation permits international transfers under specified conditions rather than requiring blanket EU-only processing; applicable law and contractual commitments must be translated into explicit placement rules. Regional replicas, edge caches, and cross-region sharding then trade capacity, communication, and coordination within those rules.

10.14.2 Global load balancing with regional replicas

Deploying replicas in multiple regions cuts wide-area latency only if the global router avoids sending users to a nearby region whose GPU queues are already saturated or whose data-residency rules forbid the request. In figure 10.24, a global load balancer selects among regions that satisfy placement policy, serve a compatible artifact version, and meet readiness and capacity requirements. Measured network latency helps choose among those eligible destinations, because geographic proximity alone does not determine completion time, while a shared registry supplies versioned artifacts to the independent regional replicas.

Replication also creates a version-management obligation. Push-based synchronization sends artifacts to each region, while pull-based synchronization lets regions fetch them after notification or polling. Neither mechanism alone guarantees consistent serving versions. Artifact verification, readiness gates, and coordinated activation determine when a region may serve the new version; hybrid synchronization combines push notification with pull verification.

During a rollout, regions may serve different versions. The application must decide whether that is acceptable and whether a request or session needs version pinning. A region-by-region canary deployment exposes the new version to one region's traffic share at a time, and each expansion passes its own error-rate, latency, and quality gate with compatible prior artifacts retained for rollback.

↪ Staged rollout and the canary observation window, sized from the required sample count and the canary's traffic share, are developed in section 12.5.5.2 (equation 12.7).

10.14.3 Edge caching with central inference

Standing up a full GPU cluster in every continent is an expensive way to reduce round-trip latency when a large fraction of distant requests ask repeated FAQ or autocomplete questions. Caching responses at the edge avoids both the wide-area round trip and repeated central inference whenever a stored response remains valid for the incoming request. Before returning a hit, the edge must enforce authorization and match the relevant model version, conversation or retrieval context,

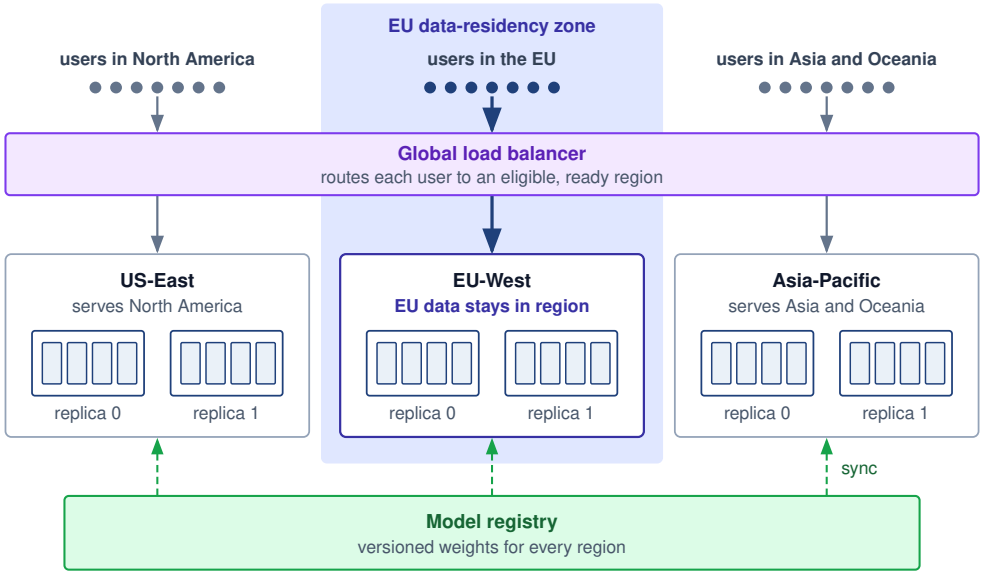

Figure 10.24: Global Load Balancing across Inference Regions: A global load balancer routes each user to an eligible region with suitable latency and ready capacity (table 10.23 gives the round-trip times this avoids). Each region (US-East, EU-West, Asia-Pacific) runs independent replicas of the same model, and EU requests can stay in the EU to meet data-residency rules. A shared model registry supplies versioned weights to every region. If a region fails, the load balancer reroutes its users only to destinations that still satisfy placement, readiness, and capacity requirements.

generation policy, and freshness requirements. Identical visible queries do not necessarily permit the same answer, and fresh-sample semantics may prohibit reuse. A hit returns within an illustrative 10 ms budget, while a miss reaches the central cluster (figure 10.25); only responses eligible for reuse should enter the edge store.

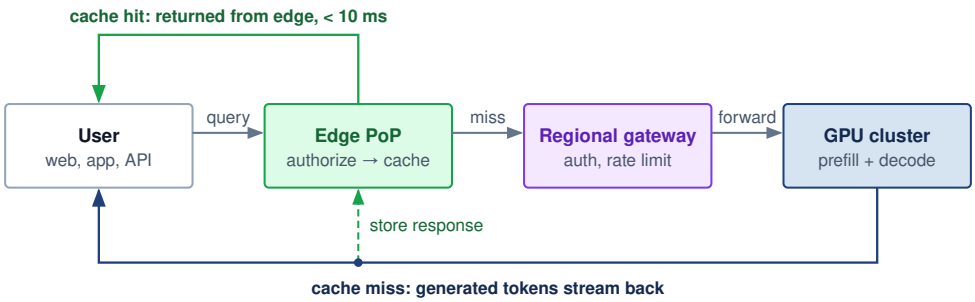

Figure 10.25: Cache Hit and Cache Miss Paths at the Edge: A query first reaches the nearest edge point of presence (PoP). After enforcing authorization and reuse eligibility, the edge targets a cache-hit response in under 10 ms without touching a GPU. On a miss, the request passes through the regional gateway to the central GPU cluster, which runs prefill and decode. Generated tokens stream back to the user, and a response eligible for reuse can be stored at the edge for later queries. The saving depends on the rate of requests eligible for reuse; open-ended generation may offer few such requests.

On the two paths of figure 10.25, cache hits avoid the backend round trip, while misses incur it. Effectiveness depends on request repeatability across workloads such as autocomplete, FAQ chatbots, open-ended chat, and code generation. Table 10.24 supplies illustrative hit-rate envelopes for sensitivity analysis; an actual design must derive hit rate and correctness from production traces.

Table 10.24: Illustrative Edge-Cache Hit-Rate Envelopes: Assumed cache-hit ranges for four workload shapes. These values are scenario inputs, not universal workload statistics. Trace-derived repeatability, response validity, and latency saved determine whether caching is worthwhile.

Workload	Cache hit rate	Suitability
Autocomplete	60–80%	Excellent
FAQ chatbot	40–60%	Good
Open-ended chat	5–15%	Poor
Code generation	20–40%	Moderate

Semantic caching[8] (caching based on embedding similarity rather than exact match) can improve hit rates for open-ended workloads. Any cache miss still requires full model execution, raising the question of whether a large model can be split across regions or must run inside a single data center.

[8] **Semantic caching**: A similarity lookup can reuse responses beyond exact string matches, but it adds retrieval work and can return an invalid answer when superficially similar queries require different responses. The similarity threshold therefore trades cache coverage against correctness.

10.14.4 Cross-region model sharding

Splitting a single model replica across data centers puts wide-area network latency directly on the request's execution path: when pipeline stages reside in separate regions, every stage crossing pays a 75–150 ms inter-region transfer against only 1–5 ms of local stage computation. That illustrative gap can slow inference by 10–100×, though the exact penalty depends on the model's stage layout. Such placement needs a reason strong enough to justify that cost, such as unavailable colocated capacity, and must be evaluated against the actual stage layout and deadline. Because cross-region sharding is so costly on the steady-state path, most production designs keep inference co-located and use other regions for resilience. The cross-region question therefore shifts from splitting one request across regions to deciding how quickly traffic can move when a region fails.

10.14.5 Cross-region failover

When a region becomes unavailable, failover must choose a healthy destination permitted to process the request and able to absorb its work. A simplified fallback tries the nearest eligible region first and, on timeout or rejection, the next one (listing 10.3). A production policy must also preserve the request's remaining deadline, version requirements, and data-placement constraints when selecting the secondary region.

Listing 10.3: Active-Active Failover: A bounded retry policy selects eligible regions and limits each attempt to the remaining request deadline.

```
# Eligibility includes placement, version, readiness, and capacity.
def route_request(user_region, request):
    attempted = set()
    for attempt_budget in (2.0, 5.0):
        remaining = request.deadline - monotonic_time()
        regions = eligible_regions(request, excluding=attempted)
        if remaining <= 0 or not regions:
            break
        region = nearest_region(user_region, regions)
        attempted.add(region)
        try:
            # The endpoint atomically admits work and enforces the deadline.
            return call_region(region, request,
                               timeout=min(attempt_budget, remaining))
        except (Timeout, RegionUnavailable, AdmissionRejected):
            # Cancel abandoned work; retries require replay-safe semantics.
            cancel_attempt(region, request)
    return unavailable_or_deadline_exceeded(request)
```

Rerouting is only part of regional failover. The destination needs accessible, authoritative conversation history to reconstruct lost KV state and enough headroom for both reconstruction and continuing traffic. If the required history is unavailable or cannot be processed there under the placement policy, a healthy replica alone cannot restore the session.

↩ Regional headroom, session rebuilding, and the gradual return of traffic to a recovered region are established in section 7.11.2.2.

10.14.6 Cost optimization across regions

Running every batch job and latency-tolerant query in the primary interactive region overpays whenever another compliant region has cheaper spot or off-peak GPU capacity. Because GPU pricing varies by provider, contract, region, and time, table 10.25 uses a hypothetical price snapshot to show how placement can balance cost against latency requirements; it must not be used for procurement.

Table 10.25: Hypothetical Regional Price and Latency Snapshot: Scenario prices and RTTs illustrate a placement calculation. Both inputs are time- and provider-specific; a deployment must refresh them from its contract and measured network paths before making a routing decision.

Region	H100 spot price	On-demand	Latency to US users
US-East	$2.50/hr	$4.00/hr	10–50 ms
US-West	$2.30/hr	$3.80/hr	30–70 ms
EU-West	$2.80/hr	$4.20/hr	75–100 ms

Cost-aware routing compares eligible regions for latency-tolerant work (listing 10.4). GPU price alone is insufficient: transfer charges, startup and idle capacity, retries, and achieved useful throughput determine the cost of completing the workload. Deadline and data-placement constraints must filter destinations before that comparison.

Listing 10.4: Cost-Aware Routing: The router filters eligible destinations before comparing completion cost for low-priority work or proximity for interactive work.

```
def route_batch_request(request):
    # Filter by placement, version, readiness, capacity, and deadline.
    regions = eligible_regions(request)
    if not regions:
        return no_feasible_destination(request)
    if request.priority == "low":
        return lowest_completion_cost_region(request, regions)
    return nearest_region(request.user_location, regions)
```

Time-of-day routing can reduce batch cost when regional price differences exceed data-transfer and coordination costs. Savings require a workload-specific calculation, and interactive traffic must retain its own latency budget. Regional placement completes the serving hierarchy, from one request's batch slot to a fleet that spans continents. The design cases that follow show how these techniques combine when a service must meet specific cost and latency targets.

10.15 Case Studies

Individual serving optimizations such as quantization, continuous batching, sharding, or caching can look decisive in isolation yet still leave a production endpoint missing its latency, memory, or cost contract when combined under real traffic. The five design cases below are illustrative composites rather than descriptions of named companies' internal deployments, and their request rates, latency budgets, hardware allocations, and update cadences follow the chapter's scenario convention. Each case binds on a different constraint: one replica's decode step, embedding scale, variable-length generation, cascade cost, or multimodal freshness. Orca and vLLM supply the previously cited scheduling and memory-management mechanisms; the scenarios combine those ideas with precision, routing, sharding, and admission control.

10.15.1 One replica against its token-latency target

Even before a serving cluster worries about routing or multi-region failover, a single replica can fail its per-token latency target on a batch of one, which leaves no headroom to batch additional requests. The first case stays inside one replica, the reference 70B model on eight NVLink-connected H100s with FP16 weights split eight ways, serving an interactive endpoint whose illustrative TPOT target is 50 ms. A batching policy can only spend the headroom that a batch-one decode step leaves, because every admitted sequence lengthens the step (section 10.6.2), so the first question is whether one batch-one step fits inside the target at all. This constructed example illustrates budget accounting, not benchmark results. Every interval is an exposed, nonoverlapping contribution to the same request's critical path, and time to first token, output quality, and throughput under load each need their own measurements.

The deployment's FP16 weights occupy 140 GB, about 17.5 GB per GPU. The assumed baseline decode step is 85 ms, about 16× the 5.2 ms that streaming each shard at full HBM bandwidth would take, so the replica runs at an MBU of about 6.1 percent (section 9.2.2). That gap, not the weight-streaming floor, is what eager dispatch, unfused kernels, and exposed synchronization must explain, and table 10.26 apportions the whole step among those causes on one absolute-time boundary. In an actual trace, overlapping work must not be counted twice when constructing such a budget.

Table 10.26: Constructed Decode Budget for the Reference Replica: Assumed nonoverlapping critical-path intervals for one batch-one decode step. Each column sums to its step duration; none is an empirical performance claim.

Exposed interval (ms)	Baseline	After precision	After fusion and replay
GEMM	34	18.7	18.7
Attention	20.4	20.4	15.3
Element-wise operations	10.2	10.2	3.4

Exposed interval (ms)	Baseline	After precision	After fusion and replay
AllReduce	10.2	10.2	10.2
Launch and runtime overhead	10.2	10.2	2
Total	**85**	**69.7**	**49.6**

The first experiment targets weight-read traffic with INT4 weight-only quantization, the bandwidth-bound decode row of table 10.11. Assume that quality validation succeeds and that exposed GEMM time falls by 45 percent, with every other interval unchanged. The step takes 69.7 ms, still above the 50 ms target, because a large cut to one interval removes only that interval's share of the step. INT8 KV-cache compression is a separate candidate whose benefit is capacity, as notebook 10.7 computes; it needs its own quality check and a new profile, so this column grants it no latency benefit.

The second experiment targets exposed attention, element-wise work, and launch gaps, using the single-replica levers of section 9.7.3. Assume a validated fused decode-attention kernel, the single-pass kernel over the KV cache of section 10.8.1, reduces attention by 25 percent, fusion reduces element-wise work to 3.4 ms, and graph replay reduces launch and runtime overhead to 2 ms. These are scenario assumptions, and each intervention would be tested separately before the combination is retained. The step becomes 49.6 ms, inside the target with almost no slack and still about 9× above the bandwidth floor.

Communication remains 10.2 ms, unchanged in absolute time even as its share of the shorter step grows, and it costs far more than launch and transfer alone predict. Spread across the 160 AllReduces of an 80-layer step, that interval implies about 64 µs per collective, against roughly 5 µs of launch and transfer, or about 0.8 ms per step. The difference is the protocol and synchronization cost of small collectives that section 9.7.2 flags for measurement, and a trace that attributed it would make communication the next experiment.

The third experiment changes how much output each step yields. Assume a draft model proposes $K = 5$ tokens in 4 ms, verification costs one round-two step, and each candidate is accepted independently with probability $\alpha_{\text{acc}} = 0.78$. Equation 10.6 gives 3.52 tokens per round, so a 53.6 ms round yields 15.2 ms per emitted token, an amortized 3.26-fold improvement in TPOT. The tokens, however, arrive in blocks. The gap between blocks is a whole round, 53.6 ms, which exceeds the same 50 ms if the endpoint also states it as a per-gap TBT target. Speculation thus meets an average-cadence contract with room to spare while missing a per-gap one on every round, which is why the TPOT and TBT targets of section 10.2 are stated separately. The gain also belongs to batch one; as the speculative-decoding policy in section 10.7.6 warns, larger batches consume the spare compute that made verification cheap.

The budget exposes bottleneck displacement without crediting a system improvement to an isolated kernel speedup. Precision changes one interval, fusion and replay change others, and speculation changes the output per round; their benefits compose only on a shared accounting boundary, with quality preserved. What the sequence produces is a set of candidate replica configurations, not a serving result. Accepting one for production means replacing the illustrative intervals with measurements, preserving the request-length distribution and offered load, and measuring SLO goodput (section 10.3) at the batch the scheduler will actually admit. Memory freed by INT4 weights can permit that larger batch, but neither the freed capacity nor a faster batch-one step shows that the service meets every target at once. Where this first case stays inside one dense LLM replica fighting per-token latency, the next case shifts to a fleet where sparse embedding lookups across terabytes of state dominate the budget before dense GPU execution even begins.

10.15.2 Embedding-heavy recommendation serving

An embedding-heavy recommendation service binds on sparse lookup rather than dense forward-pass compute. Modeled on feeds, advertising, and content ranking, it processes billions of requests per day under a p99 budget below 10 ms for the ranking service, nested inside the roughly 100 ms end-to-end tail target of Archetype B (DLRM at Scale). The service spans hundreds of model

variants and trillions of unique sparse features. The serving stack separates sparse lookup from dense ranking:

```
User request → Feature collection → Embedding lookup → Model inference → Response

                    |                    |                    |
                    v                    v                    v
            Feature Store        Embedding Servers      GPU Inference
            (CPU, DRAM)        (CPU + SSD, 1000s)     (GPU, 100s)
```

The hybrid sharding strategy in lighthouse 10.1 supplies the sparse path: replication protects hot entries, column sharding spreads warm-entry access, and row sharding places the cold tail. Feature-level batching then combines lookups across requests. The remaining design decision is where each stage runs. In the GPU-CPU hybrid architecture, memory-rich CPU servers serve sparse embeddings while GPUs execute dense ranking. Table 10.27 connects that placement to the end-to-end latency budget.

Table 10.27: Illustrative Hybrid CPU/GPU Architecture: Illustrative hardware-workload mapping for the recommendation scenario. CPU + SSD handles the bandwidth-bound sparse path; CPU handles low-arithmetic-intensity feature processing; GPU handles the compute-dense ranking forward pass. The component budgets exclude queueing and response delivery; their sum alone does not establish the p99 target.

Component	Hardware	Latency	Throughput
Embedding lookup	CPU + SSD	2 ms	50M lookups/s
Feature processing	CPU	1 ms	10M ops/s
Dense ranking	GPU	1.5 ms	100K infs/s

The serving lesson is that recommendation latency is often dominated by embedding lookup rather than dense model inference. Feature-parallel batching and a hybrid CPU-GPU architecture match the hardware to the sparse and dense halves of the workload, which sets up a contrast with GPT-style API serving where request variance and autoregressive state dominate.

10.15.3 A multi-model generation API

A multi-model generation API binds on request variance, because short prompts, long-context requests, and different model sizes share capacity while users expect predictable TTFT. The platform serves millions of requests per hour under a time-to-first-token target below 2 s, with models from billions to hundreds of billions of parameters and contexts up to 128K tokens. The serving stack separates routing, scheduling, cache management, and shard groups:

```
API Gateway → Rate Limiting → Request Router → Model Cluster → Response Streaming

                    |
                    v
            +-------------+
            | Model Pool  |
            | +---------+ |
            | | Large   | |
            | | 8×H100  | |
            | +---------+ |
            | +---------+ |
            | | Smaller | |
            | | 4×A100  | |
            | +---------+ |
            +-------------+
```

The main design decision is to prevent long prompts from blocking decode service for everyone else. An Orca-style design uses continuous batching to refill completed slots. Chunked prefill limits interference from long prompts by interleaving prompt chunks with ongoing generation. The TTFT target must be validated across prompt lengths; admitting a 128K-token prompt does not imply that it meets the short-request latency target. Slot reuse and long-prompt interference are separate concerns. Continuous batching addresses the first and chunking bounds the second, so neither mechanism alone establishes the utilization or tail latency the API will see.

The tensor-parallel degree must satisfy weight and KV-cache capacity while leaving enough latency budget for communication. This scenario allocates eight H100 GPUs per large-model shard group connected via high-speed NVLink for intra-node communication, with consistent hashing at the gateway to maintain session affinity and enable compatible cross-turn KV cache reuse. The eight-device allocation is a design choice to validate against the model, precision, and admitted context lengths, not a requirement of GPT-class models in general.

To prevent noisy neighbors from exhausting capacity, the illustrative API applies hierarchical rate limiting across multiple granularities: per-API-key request arrival limits, per-API-key token-per-minute consumption limits, organization-level capacity quotas, and global model capacity ceilings. Within those ceilings, capacity can also shift between models on queue depth, subject to model-loading delay and each service's minimum allocation, using the autoscaling signals of section 10.13.

LLM APIs are governed by variance control. Continuous batching reuses execution slots, prefix caching avoids compatible repeated work, and multi-tier rate limiting bounds admitted demand. Their combined effect on latency still depends on prompt scheduling and available capacity. Search ranking faces a different shape of the same problem. Many cheaper models must cooperate under one strict deadline.

10.15.4 Search ranking under a deadline

The search-ranking scenario binds on cascade cost, because each query coordinates many smaller models for query understanding, document relevance, and result ranking under one end-to-end budget rather than serving one large model. The system processes billions of searches per day under a target below 200 ms while evaluating thousands of candidate documents per query, so screening candidates with cheaper models reserves the expensive ranking budget for a smaller set.

To meet this latency budget while evaluating thousands of documents, the ranking cascade in figure 10.26 screens candidate pools across four progressively expressive tiers. Early lexical and approximate-nearest-neighbor (ANN) embedding retrievers discard 99 percent of candidates within a 10 ms budget, reserving compute-intensive cross-attention transformers for the final top 100 items. This saving creates a recall constraint. A relevant document discarded upstream cannot be recovered by a more accurate downstream ranker, so each stage's candidate budget must be evaluated for retained relevant items as well as latency.

↩ The vector indexes behind the L0 retrieval stage, and the storage cost of keeping them fast enough for a serving deadline, are established in section 4.9.1.

The central design decision is progressive refinement through a ranking cascade rather than one expensive model applied to every candidate. Each stage of the cascade spends a larger latency budget on a smaller candidate set (table 10.28), so the expensive L3 transformer is affordable only because L0 through L2 have already cut the candidates by four orders of magnitude.

Table 10.28: Illustrative Ranking Cascade: Each stage evaluates fewer candidates with a more expensive model. The final stage has the largest latency allocation in this scenario; candidate reduction makes that expense affordable.

Stage	Model complexity	Candidates	Latency budget
L0 (Retrieval)	Lexical and ANN embeddings	1,000,000 → 10,000	10 ms
L1 (First pass)	Two-tower or GBDT	10,000 → 1,000	20 ms
L2 (Second pass)	Deep MLP (DLRM)	1,000 → 100	50 ms
L3 (Final rank)	Cross-attention transformer	100 → 10	100 ms

The cascade reduces final L3 evaluations by 10,000× compared to running L3 on all candidates. If the stage budgets are treated as additive model work and L3 cost scales linearly with candidate

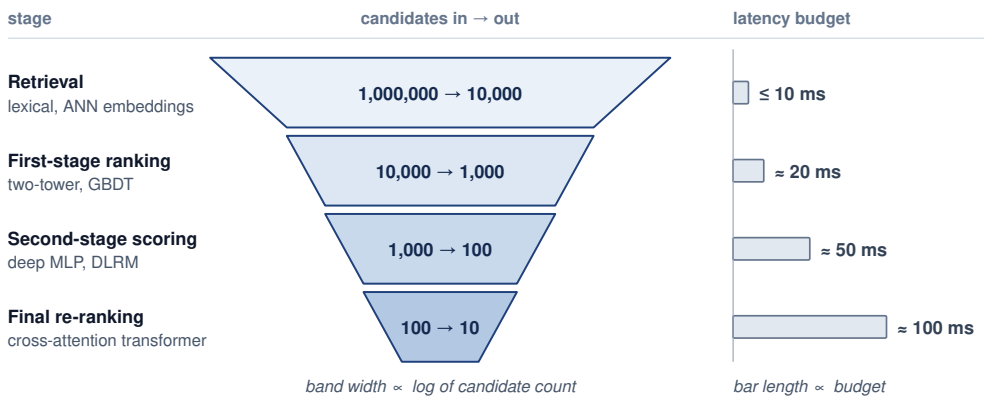

Figure 10.26: Ranking Cascade Architecture: To optimize latency and cost, search and recommendation systems use a cascade of increasingly complex models. Early stages (Retrieval) filter millions of candidates down to thousands using fast, cheap heuristics. Later stages use expensive, high-precision models only on the most promising candidates.

count, the L3 allocation implies 1 ms per candidate. Applying L3 to the full million-candidate set would then consume about 1000 s of model work, while the cascade budget sums to 180 ms. The resulting reduction is roughly 5,600× under these assumptions. This accounting estimate is not a measured latency speedup; batching and parallel execution can change elapsed time without changing the candidate count.

A deadline-constrained ranking ensemble can use speculative execution when it has a validated fallback-quality policy. Here, speculation means launching candidate ranking work in parallel under a deadline, not draft-token verification for LLM decoding. The final model starts on predicted inputs, such as cached or default scores, in place of the earlier models' outputs, and its speculative result is used only if the actual outputs miss the deadline:

```
# Instead of sequential:
#   q1 = model1(query)
#   q2 = model2(query)
#   q3 = model3(query, q1, q2)

# Speculative parallel:
async_q1 = async model1(query)
async_q2 = async model2(query)
async_q3 = async model3(query, predicted_q1, predicted_q2)

# Use actual results if they arrive in time, otherwise use speculative
```

Each sub-request carries an explicit remaining-time budget, and workers order their local queues by deadline proximity rather than arrival order. When a queued sub-request cannot complete before its budget expires, the worker returns a cached or default score immediately instead of timing out the parent query:

```
Worker queue: [Doc1: 50 ms left] [Doc2: 30 ms left] [Doc3: 80 ms left]
                                    ↑ Process first

If deadline will be missed:
  Return cached/default result rather than timing out
```

The serving lesson is that ranking cascades reduce cost by spending expensive models only on a shrinking candidate set, at the risk of losing relevant items upstream. Deadline propagation and priority scheduling expose the remaining time budget; meeting it also requires sufficient capacity and a validated fallback when work cannot finish. The next case adds a freshness constraint that deadline scheduling alone cannot solve.

10.15.5 Multimodal recommendation and freshness

The multimodal recommendation scenario binds on freshness under a tight ranking SLO: video understanding is expensive, new content arrives continuously, and user modeling remains online. The system ranks millions of videos per second under a p99 target below 50 ms while ingesting millions of new videos daily across video, audio, text, and interaction signals, which is why online and offline work must separate. The serving stack separates content and user paths:

```
User request → User embedding → Candidate videos → Video understanding → Ranking

                    |                 |                  |
                    v                 v                  v
              User Tower        Video Cache        Vision Models
            (Transformer)     (Pre-computed)      (On-demand)
```

The scenario separates user understanding (online) from content understanding (offline) through a two-tower architecture with caching.[9] The user tower runs online on every request, while the video tower refreshes in hourly offline batches (table 10.29).

[9] **Two-tower architecture**: Encodes user features and item features through separate neural networks ("towers") into embeddings, then scores via dot product. Its serving advantage is that item embeddings can be precomputed offline and cached, reducing online candidate generation to the user tower forward pass and vector similarity search. Search latency depends on index size, placement, and the required recall. The trade-off is reduced model expressiveness from late interaction, since the towers cannot attend to each other's features during encoding.

Table 10.29: Illustrative Two-Tower Update Cadence: User and video towers run on different cadences: user features change second-to-second and must run inference online; video features change much more slowly and are precomputed in hourly batches. This asymmetric scheduling is the design that takes vision inference out of the critical path for almost every request.

Tower	Update frequency	Latency	Compute
User tower	Real-time	5 ms	GPU (online)
Video tower	Hourly	N/A	GPU (batch)

Video embeddings are precomputed and cached, eliminating vision inference from the critical path for most requests. A newly uploaded video without a cached embedding needs an explicit policy: process it before admitting it to recommendations, serve a validated fallback, or run vision inference on demand if the request budget permits. The hourly refresh cadence alone does not meet the shorter ingestion targets below; those targets require a separate path for new content.

The online path below separates candidate generation from a later dense reranker. Its user-tower budget is distinct from the reranker budget:

```
User features → CPU preprocessing (1 ms)
              → Embedding lookup (2 ms, CPU+DRAM)
              → User tower (5 ms, GPU)
              → Dense reranking (10 ms, GPU)
              → Response formatting (1 ms)
```

New uploads enter an ingestion queue with priority tiers, so time-sensitive content is embedded within minutes and bulk imports absorb hours of slack, although priority only changes who waits when ingestion capacity falls short. The modality embeddings (video, audio, text) are combined late by a small fusion model, so updating one modality's encoder requires a compatibility check and, when its embedding contract changes, regeneration of the cached embeddings.

The serving lesson is that freshness requires separating online and offline work. Two-tower architectures make aggressive caching possible because item embeddings can be refreshed asynchronously, while priority-based processing spends the freshest compute where it has the largest user impact.

10.15.6 Cross-cutting observations

Trying to serve LLM token generation, sparse recommendation ranking, and multimodal retrieval through one generic model server forces every workload to pay for scheduling and memory abstractions it does not need. Across these cases, the common pattern is specialization rather than a universal serving stack. Embedding and retrieval stay separate from ranking and generation, hardware is matched to each stage, caches accept bounded staleness, candidate sets shrink before expensive work, and deadlines propagate so quality-latency trade-offs stay explicit under pressure. Table 10.30 distills the same choices along five design axes for the three workload families the cases cover. Its entries are candidate combinations for the illustrated constraint regimes, not prescriptions, because a model class alone does not fix a scheduler or topology, which also depend on model size, request mix, available hardware, and service targets.

Table 10.30: Serving Architecture Dimensions by Workload Type: The rows identify candidate design choices along five architectural dimensions. The workload's state, scheduling, and hardware constraints guide selection; the service contract determines whether a combination supplies useful capacity.

Dimension	Vision/embedding	LLM (autoregressive)	Recommendation
Batching	Static/dynamic (uniform inputs)	Continuous (variable outputs)	Feature-parallel (sharded embeddings)
Memory	Preallocated (predictable)	Paged (variable KV cache)	Distributed (embedding tables)
Scheduling	FCFS (uniform cost)	Preemptive (high variance)	Priority-aware (SLO tiers)
Topology	Single-GPU replicas	Tensor/pipeline parallel	Hybrid CPU-GPU sharding
State	Stateless	Request state; optional session cache	Stateless or session features

The case studies show why serving architectures must be adapted to their workloads. Applying conventional web-service assumptions to GPU-bound systems creates the fallacies and pitfalls examined next.

10.16 Fallacies and Pitfalls

A DevOps team that applies its standard web-server load-balancing rules to a new fleet of GPU inference nodes can see tail latency climb sharply, because those rules ignore the extreme variance of autoregressive token generation. Web services also face variable work and queueing, but an inference fleet adds model residency, accelerator scheduling, and per-sequence state to those familiar concerns. The fallacies below arise when a serving design ignores these constraints or treats one optimization as a complete solution.

Fallacy: *Serving at scale is synonymous with LLM serving.*

The misconception leads to over-focus on LLM-specific techniques. As section 10.5.1 showed, recommendation and ranking models can consume most AI inference cycles on large consumer platforms (Gupta et al. 2020; Hazelwood et al. 2018). A practitioner who knows only continuous batching and KV cache management is unprepared for the feature-parallel batching and embedding sharding that dominate that traffic. Technique selection must match the actual workload.

Pitfall: *Using training infrastructure for production serving.*

Training and serving have different requirements, as table 10.2 summarized. Interactive serving must satisfy per-request latency targets under external arrivals, so its batch size depends on the workload and waiting budget, and its recovery policy must account for in-flight requests, retries, and the availability target rather than a checkpoint. Hardware can be shared or repurposed, but a training configuration does not establish serving readiness.

Fallacy: *Continuous batching solves the whole LLM serving problem.*

Continuous batching reuses slots released by completed requests, but it addresses only one dimension of the problem. Prefill can remain a bottleneck for long contexts because scheduling subsequent decode iterations does not remove the quadratic attention work of full-attention prefill. Section 10.7 demonstrates that KV cache memory, not compute, often limits batch size. Network bandwidth between sharded model components can dominate latency for large models. The value of continuous batching depends on the workload; slot reuse alone does not resolve these other constraints.

Pitfall: *Sizing the fleet by average throughput.*

Average throughput alone does not establish capacity for bursts. With unchanged work per request and service capacity, a 25 percent traffic increase from 80 percent utilization reaches 100 percent utilization, leaving no spare capacity to absorb variance. Whether the resulting queue violates an SLO depends on burst duration, the waiting budget, and admission policy. Cluster-level cold starts can delay relief for multiple minutes. Capacity planning must therefore account for peak load, burst duration, ready headroom, and the work the service can defer or reject.

Fallacy: *Load balancing does not matter much for inference.*

Simple load balancing strategies like round-robin can distribute request counts while leaving work uneven. The static placement analysis in section 10.11 motivates two probes, but its occupancy bounds do not predict production queue lengths. Serving comparisons must preserve arrivals, request lengths, cache affinity, and replica capacity. Work-aware routing earns its place when it improves measured tail latency without excessive coordination cost.

Pitfall: *Ignoring the serving tax.*

Distributed inference adds network round-trips and shard coordination to the serialization, routing, and queueing costs that can already surround a single-machine model. This serving tax depends on deployment topology and load. In the illustrative budget of section 10.3, 70 ms of model execution leaves only 30 ms for every other stage of a 100 ms request target, and queueing alone can consume that remainder as load rises. A kernel benchmark measures only the first term of equation 10.1; the serving latency is the measured sum along the complete request path.

Fallacy: *More GPU memory always means larger batch size and higher throughput.*

Larger GPU memory raises throughput only when the extra capacity supports useful work within the latency budget. Even when memory bandwidth is saturated, a larger decode batch can increase tokens per second by sharing weight reads across more sequences. That gain stops when compute, per-sequence KV-cache traffic, or the SLO becomes binding. The bottleneck hierarchy distinguishes these limits.

For an H100 with 80 GB memory and 3.35 TB/s HBM3 bandwidth, capacity and transfer rate are separate limits. More capacity may admit additional sequences, but the resulting batch must still meet compute, bandwidth, and latency budgets. Once extra sequences no longer improve useful throughput within those budgets, adding memory alone cannot help. Capacity planning therefore measures the bottleneck at the intended batch size and context length rather than assigning one permanent bottleneck to the model.

Pitfall: *Neglecting multi-tenancy isolation until production.*

In development and staging, single-tenant deployments work well. In production, noisy neighbors cause sudden, unpredictable performance degradation that is difficult to diagnose and resolve. A tenant bursting to $5\times$ normal traffic can degrade latency for all other tenants on shared infrastructure. As emphasized in section 10.12, resource quotas, priority scheduling, and bulkhead isolation must be designed into the system from the start, not retrofitted after production incidents.

Fallacy: *The lowest supported precision is always the best precision.*

Aggressive quantization (INT4 weights, INT4 KV cache, FP8 activations) can degrade model quality in ways that standard benchmarks miss but users see. Perplexity on a held-out dataset may change by less than 1 percent while responses degrade for edge cases, rare languages, or complex reasoning. A serving system that doubles its tokens per second with a format that lowers answer quality has not doubled its useful capacity, because throughput and quality form a frontier and a precision choice is a point on it. The disciplined approach is targeted quantization: measure per-layer sensitivity under representative workloads, restrict lower bit widths to the tensors that tolerate them, and evaluate KV caches, endpoint layers, and attention logits empirically rather than by a presumed sensitivity ranking. Calibration on representative data, followed by evaluation on diverse quality benchmarks reported beside the latency and throughput results, comes before any quantized model reaches production (section 10.9).

Pitfall: *Treating MoE expert count as a free scaling knob.*

Adding experts to an MoE model raises total parameters without proportionally raising active per-token arithmetic, which makes expert count look free. Every expert still occupies memory whether or not a token selects it, as section 10.10.4.1 prices for resident versus active parameters, so more experts mean more devices or less KV headroom. A distributed placement also adds dispatch and combine communication to every layer. The realized cost depends on experts selected per token, token-to-expert skew, expert capacity, topology, and batching, and no universal expert count marks where the design starts to lose. The change is justified only when a matched-quality comparison shows that active-compute savings exceed routing, imbalance, memory, and communication costs over the intended workload envelope.

These fallacies and pitfalls identify what a serving evaluation must test: traffic variance, cost per completed request within the SLO, output quality beside throughput, and interference across tenants. Correct technique selection must be followed by measurement under the intended load and failure conditions.

10.17 Summary

Serving at scale turns trained models and distributed hardware into a service under continuously arriving work. Training optimizes aggregate progress over a run; serving must also meet user-visible latency targets while paying recurring operating costs. The useful capacity of the fleet is the work it completes within its service contract, not the work its accelerators execute in isolation. Test-time compute sharpens the point. When quality is bought by spending more compute per query, the scaling curve that drove the fleet's growth moves into the request path, and every increment of quality is paid again on every request.

Serving performance depends on the balance between compute and memory bandwidth. In the terms of the single-machine iron law recalled in section 1.5.4, long-prompt prefill can be compute-bound (O/R_{peak}), while small-batch decode is often bandwidth-bound (D_{vol}/BW). Each decode iteration reads active weights and the sequences' KV state; batching shares weight reads across more tokens, whereas longer contexts increase KV traffic. Single-replica efficiency (η_{hw}) therefore depends on phase, batch size, and context length. Variable output lengths make fixed batch membership waste capacity after shorter requests finish.

To maximize throughput under memory-bandwidth constraints, modern serving architectures co-design the memory hierarchy and the iteration scheduler. Continuous batching recovers idle slots by inserting new requests at token boundaries. PagedAttention reduces physical memory fragmentation through fixed-block allocation while retaining partial-page waste. Prefill/decode disaggregation assigns prompt processing and token generation to separate hardware pools sized around their measured compute, memory, and communication costs.

When model state exceeds device memory, sharding spreads a replica across accelerators; when demand exceeds a replica's service capacity, additional replicas supply aggregate throughput. In Fleet Law terms, sharding divides the compute term T_{compute}/N of equation 10.12 but adds communication and synchronization $(T_{\text{comm}}(N), T_{\text{sync}}(N))$ governed by interconnect latency floors (α) and transfer bandwidth (β). Replication, routing, isolation, autoscaling, and multi-region placement add coordination costs and wide-area round trips to the same budget, and their tail

control requires headroom. Across the serving hierarchy, the engineering objective is to increase useful completed work while preserving the latency and reliability contract under load.

Key Takeaways: Serving capacity is constrained by its contract

- **Serving cost recurs throughout deployment**: Training is paid per model version, while inference expense accrues with traffic (principle 17). In the chapter's high-volume recommendation scenario, lifetime serving cost reaches 788.4× the training cost, making efficiency a recurring lever, measured as cost per completed request at a fixed request mix, quality, and SLO.
- **SLO goodput is the capacity measure**: A serving fleet's capacity is the rate of completed requests that meet their TTFT and per-token targets for a fixed request mix, so batching, sharding, caching, and autoscaling help only when they raise that rate. Aggregate throughput overstates capacity whenever queueing, first-token delay, or slow token delivery makes completed work miss its contract.
- **Batching trades waiting for reuse**: Throughput saturates past a knee set by the ratio of fixed to per-request batch cost, and a batching window caps formation delay but not queueing for a busy replica. Little's law then sizes the concurrency and replica count that a target rate and mean latency imply.
- **Decode turns memory into capacity**: Long-prompt prefill often binds on compute, while small-batch decode rereads active weights and expands private KV state token by token (principle 18). Continuous batching, PagedAttention, prefix reuse, and KV compression matter because they convert scarce HBM bandwidth and memory into admitted requests.
- **Prefill and decode compete for a shared worker**: A monolithic long-prompt prefill lands in every active stream's token gap. Chunked prefill bounds that gap by moving the cost into the long prompt's own time to first token, while separate prefill and decode pools remove the interference only when their KV-cache transfer is repaid in SLO goodput.
- **Model class sets the primitive**: Vision models benefit from predictable static or dynamic batches, recommenders from feature-parallel embedding sharding, and LLMs from iteration-level scheduling. A fleet scheduler must match the workload shape instead of applying one universal batch-size rule.
- **Distribution always sends a bill**: Sharding, routing, isolation, and failover buy capacity or resilience by spending communication and coordination. Tensor parallelism pays collectives on every generated token, so it is normally kept inside one fast interconnect domain. Power-of-two choices (principle 19) needs only two probes, which for variable-length generation should compare remaining work, with atomic admit-or-reject at the destination because probe results go stale.

The serving cost dominance principle (17) makes local efficiency valuable only when the service can turn it into useful capacity. Memory released by quantization may admit more sequences, but those sequences still compete for execution time and can push the endpoint beyond its latency budget. The economic comparison therefore holds the request mix, quality requirement, and service targets fixed before asking whether a change lowers the cost of completed work.

What's Next: From data center to edge

The serving designs developed here coordinate accelerator replicas under data-center memory, network, and availability assumptions; even the global stack's edge points of presence only answer from cache and leave model execution in the clusters. Many applications also run models on smartphones and other personal devices, where sustained power is a few watts, memory is shared with foreground work, and connectivity is intermittent. When those devices must also adapt to local data, the fleet has to decide where learning happens, which device

updates it accepts, and how a population improves a shared model without collecting raw examples.

↪ On-device adaptation under device resource budgets and federated learning across device populations are developed in chapter 11.

11

Edge Intelligence

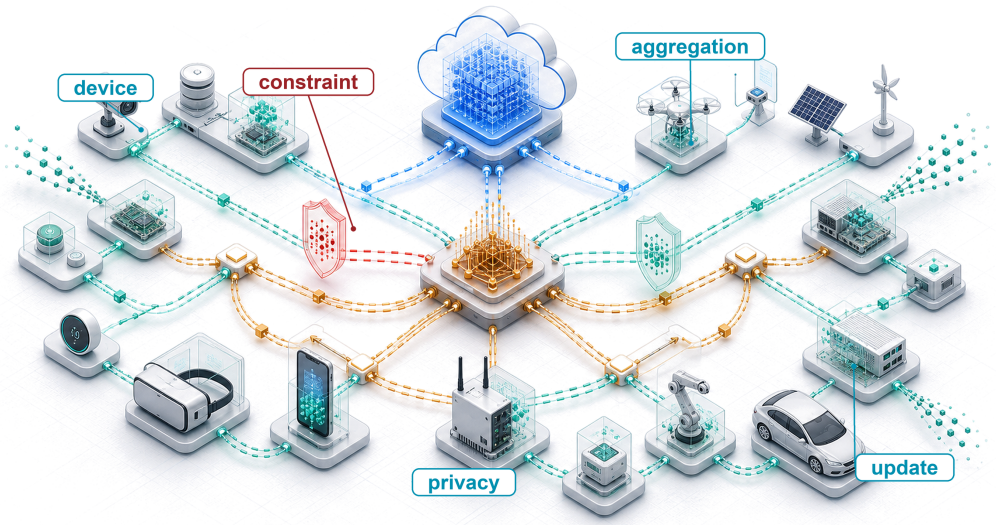

11.1 The Edge Learning Paradigm
11.2 Why Learn on the Device
11.3 Design Constraints
11.4 Model Adaptation
11.5 Data Efficiency
11.6 Federated Learning Algorithms
11.7 Federated Learning Systems at Scale
11.8 Continual Adaptation Under Edge Budgets
11.9 Integrating the Edge Intelligence Pillars
11.10 Production Integration
11.11 Engineering Challenges and Mitigations
11.12 Fallacies and Pitfalls
11.13 Summary

Purpose

When does learning need to move from the data center to the devices it serves?

A model trained in a data center encounters new conditions when deployed to a user's device. Vocabulary, lighting, acoustics, and usage patterns may differ from the distributions that shaped its parameters, and the gap can widen as behavior evolves. Centralized training cannot anticipate every context, while privacy constraints or limited connectivity may prevent local observations from reaching it. Learning where the data is produced can improve local behavior, but running predictions offline does not by itself justify training there. Adaptation must provide enough benefit to justify its memory, battery, and thermal costs without disrupting the device's primary work. A training job that fits on an idle phone may become unacceptable when the user starts another application. Even successful local learning leaves a fleet problem: improvements must be evaluated across devices whose capabilities, data, and availability differ. A policy that accepts updates only from well-connected devices may improve measured results while neglecting users whose devices rarely qualify. Making one device learn efficiently therefore does not establish that a population can learn reliably or that exchanging model updates protects private information. The engineering decision joins the value of local adaptation with the cost of participation and the evidence needed to trust shared improvements. In C^3 terms, compute must fit each device, communication must carry useful updates within limited connectivity windows, and coordination must preserve a representative learning population despite dropout and unequal resources.

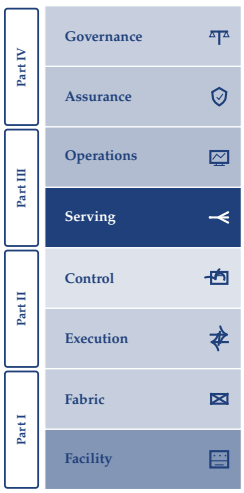

Learning Objectives

- Contrast centralized cloud training, on-device learning, and federated learning by analyzing their data flow, privacy properties, and operational trade-offs
- Quantify edge-device training overhead relative to inference-only deployment across memory, compute, and energy
- Derive an admission condition for background adaptation from energy, thermal, and foreground-latency slack
- Select model-adaptation and data-efficiency methods by the memory term each shrinks, its forgetting risk, and its storage cost
- Analyze how client sampling, dropout, and non-IID data determine which population a federated model learns from and its time to target
- Design federated learning systems that aggregate privacy-preserving updates within uplink, secure-aggregation, and version-compatibility limits
- Explain the operational controls needed for on-device learning: device-aware deployment, distributed validation, privacy-preserving monitoring, and rollback across heterogeneous populations

11.1 The Edge Learning Paradigm

IMAGINE a voice assistant trained on millions of generic audio samples in a cloud data center. When deployed, it may struggle with a regional accent or household vocabulary. Without local adaptation, its performance can decline as usage diverges from the training distribution. **Edge intelligence** places inference, adaptation, and coordination on or near the devices that observe the data, allowing models to respond to local context under power, memory, privacy, and connectivity constraints.

The serving fleet already has to meet latency targets despite failures and uneven demand. At the edge, the operator also loses control over when a device has power, connectivity, or spare memory. The unit of design remains the fleet, but its participants can disappear between selection and completion. Local adaptation matters here because its cost determines which users can contribute to population learning.

↩ Serving fleets that hold latency service-level objectives (SLOs) through replica failures and uneven demand are established in chapter 10.

Definition 11.1: On-device learning

On-device learning is the local training or adaptation of machine learning models directly on deployed hardware, using data that remains on the device; a local update needs no server connection, although sharing its result does.

1. **Significance**: It enables personalization and continued adaptation where connectivity is intermittent or raw data may not leave the device. In single-machine iron-law terms, each update adds backward and optimizer work to the compute term $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$ that foreground inference also needs, and it draws on a separate battery and thermal budget that the admission test of equation 11.2 later makes explicit.
2. **Distinction**: Unlike edge inference, which only executes a fixed model, on-device learning runs a local optimization loop (forward and backward passes plus parameter updates) that requires more memory and compute than the forward pass alone.
3. **Common pitfall**: A frequent misconception is that on-device learning requires training from scratch. In reality, it is almost always **on-device adaptation**, which fine-tunes a small subset of a pretrained base model's parameters to fit the data distribution of a local environment.

Consider one federated round for the voice assistant. A coordinator selects eligible devices and distributes a shared model version and a bounded training policy. Each selected device trains on local examples when its policy admits background work, then returns a model update tagged

with the starting version. Some devices lose connectivity or exhaust their resource allowance. The coordinator closes the round according to its deadline and participation requirements, combines compatible updates, evaluates the candidate shared model, and publishes an accepted version for later download. Local personalization can remain on the device instead of entering that aggregate.

This round has the shape of Archetype C (Federated MobileNet), with phone-class clients, private local data, and intermittent availability, and it exposes the C^3 trade-off before any adaptation method is chosen. The Fleet Law (equation 1.3) gives it a time budget, with one change of reading. In a data center the compute term divides one job's work across N devices. Here every selected device trains its own copy on its own data, so the term is set by local training on the slowest device the round still waits for rather than divided by N. The communication term is the update transfer over a wide-area uplink. The synchronization term is the wait for enough participants to report. The overlap term credits whatever downloading and training proceed concurrently.

More local compute can reduce the frequency of communication, but devices then need longer uninterrupted training windows and their models can drift toward different local distributions. Short deadlines improve coordination speed while excluding slow participants. The fleet must therefore account for whose updates complete, rather than merely how many devices were selected. Device budgets, adaptation methods, and data reuse determine the work inside the round; aggregation and scheduling determine whether that work improves the shared model. A round of this kind, repeated across a population, is federated learning.

↩ Archetype C (Federated MobileNet), the phone-class federated workload with private local data and intermittent availability, is established in section 1.6.1.3.

Definition 11.2: Federated learning

Federated learning is a decentralized training paradigm where distributed devices collaboratively train a shared model using local data while exchanging only model updates (gradients or weights).

1. **Significance**: It transforms the constraint of data locality into a privacy feature. Within the Fleet Law (equation 1.3), federated learning is constrained by wide-area bandwidth (BW) in its communication term and by the extreme heterogeneity of the fleet, where device-specific efficiency (η_{hw}) and availability can vary by orders of magnitude.
2. **Distinction**: Unlike centralized training, where data is moved to the compute, federated learning moves compute to the data. Raw training examples can remain local, but model updates, participation metadata, and protocol messages still cross the device boundary and may carry personal information.
3. **Common pitfall**: A frequent misconception is that federated learning is "inherently private." In reality, model updates themselves can leak information through gradient inversion attacks, which search for inputs whose gradient matches a submitted update and so reconstruct training examples, requiring additional protections like differential privacy or secure aggregation.

Keeping raw examples local reduces one source of exposure, but an update can still reveal information about the examples that produced it. Secure aggregation lets the coordinator obtain a combined update without inspecting each client's contribution under the protocol's trust assumptions. Differential privacy[1] bounds how much a specified user's or record's participation can change the released result, usually by limiting contributions and adding calibrated noise. Neither protection follows automatically from federation. Their costs and assumptions belong in the round design alongside dropout handling and communication.

A smartphone predicting user text input, a smart home device responding to household routines, and an autonomous vehicle interpreting local driving conditions can all encounter variation that centralized training data underrepresents. The smartphone meets user-specific phrasing. The home device observes seasonal shifts in household routines. The vehicle faces local road conditions, weather patterns, and traffic habits that differ from its training environment. Such gaps make a case for on-device learning (definition 11.1), moving machine learning from centralized training to distributed learning across heterogeneous devices that operate under tight power, memory, and thermal limits. Because local training spends those scarce device budgets, a distribution gap is a

[1] **Differential privacy**: A mathematical framework that bounds information leakage by adding calibrated noise to computations. In edge learning, that noise becomes a systems cost because noisier updates usually require more rounds or more data to reach the same utility.

↪ Differential privacy definitions, noise mechanisms, and privacy accounting are developed in section 13.9.

case for adaptation rather than a verdict; an engineering analysis must first establish when local learning is necessary and whether a given update measurably improves behavior.

11.2 Why Learn on the Device

A distribution gap justifies spending device resources only when centralized retraining cannot close it in time or within policy. A cloud-trained model that relies on a server round trip or a raw-data upload for every adjustment fails whenever a user has an unfamiliar accent, a remote sensor loses connectivity, a regulation forbids exporting personal traces, or millions of devices would overwhelm backend ingestion. Those four pressures (personalization, latency and availability, privacy, and infrastructure efficiency) motivate shifting learning to the edge (T. Li et al. 2020).

Personalization matters when deployed models encounter usage patterns and data distributions that their training environments represent poorly. Local adaptation can refine behavior using user-specific data, including linguistic preferences, physiological baselines, sensor characteristics, or environmental conditions. High inter-user variability justifies changing weights only if the adapted model measurably beats a fixed model with local context.

Latency and availability constraints determine where an update can run and how soon it can become useful. In edge computing scenarios, connectivity to centralized infrastructure may be unreliable, delayed, or intentionally limited to preserve bandwidth or reduce energy consumption. When adaptation is justified, local learning permits updates during disconnected periods. The update must still fit the device's resource budget, and its completion deadline is separate from the response deadline for foreground inference.

Privacy is the third driver. Many applications involve sensitive or regulated data including biometric measurements, typed input, location traces, or health information. Local learning can reduce privacy exposure by keeping raw data on the device and limiting transmission to centralized systems. This can simplify compliance engineering, but it does not by itself establish adherence to regulations such as the General Data Protection Regulation (GDPR), the Health Insurance Portability and Accountability Act (HIPAA) (U.S. Department of Health and Human Services 2026), or region-specific data sovereignty laws.

Infrastructure efficiency is the economic driver. Centralized training pipelines require backend infrastructure to collect, store, and process user data from potentially millions of devices. Shifting learning to the edge distributes computation and can reduce raw-data ingestion. The resulting savings must be compared with device energy, repeated model transfers, and coordination overhead; whether moving work off the server lowers end-to-end cost depends on that whole sum, not on the server bill alone.

11.2.1 Simpler alternatives first

Before paying that device and coordination cost, an engineering team should test whether a static model paired with local state already meets the requirement. Shipping an on-device gradient-training loop when a local feature store or lookup table suffices spends memory, battery, and rollback margin for no measurable gain, whereas simpler alternatives often achieve comparable quality with far lower operational overhead. Three such alternatives often satisfy personalization and adaptation requirements without the complexity of local training (table 11.1).

Table 11.1: Alternatives to On-Device Learning: Simpler personalization mechanisms can satisfy many adaptation requirements without local gradient updates, so on-device learning should be reserved for cases where privacy, latency, connectivity, or measured quality requirements rule out these options.

Alternative	Mechanism	When it fits
Feature-based personalization	Store user preferences, interaction history, and behavioral features locally, then feed those features into a static model rather than adapting model weights.	News recommendation systems can store topic preferences and reading patterns locally, then combine those features with a centralized content model.

Alternative	Mechanism	When it fits
Cloud-based fine-tuning with privacy controls	Fine-tune centrally on data authorized for transfer, with access controls and, where required, a specified differential-privacy mechanism. Federated analytics and secure aggregation (Bonawitz et al. 2017) can support aggregate statistics but do not themselves provide centralized training examples.	Centralized updates fit when authorized data transfer, update latency, and the specified privacy guarantees meet the application requirements.
User-specific lookup tables	Combine global models with personalized retrieval mechanisms by maintaining a lightweight table for frequently accessed local patterns.	Personalization benefits are needed with minimal computational and storage overhead.

The decision to implement on-device learning should follow evidence that simpler alternatives cannot meet the deployment requirements, such as restrictions on transferring training data or unreliable connectivity when adaptation is needed. Latency is the requirement most often mistaken for such evidence. For applications whose foreground deadlines are shorter than the network round-trip time, cloud inference is architecturally infeasible. Camera processing under 33 ms, voice response under 500 ms, AR/VR motion-to-photon latency under 20 ms, or safety-critical control under 10 ms all face network round-trip times typically ranging from 50 to 200 ms. When the network cannot meet the foreground deadline, inference must remain local. A fixed model also executes offline, however, so the deadline settles where the model runs and leaves open where it learns. Local training earns its engineering investment only when a measured adaptation gain justifies its added resource and operational costs.

11.2.2 Where adaptation pays

Every application domain must justify adaptation against its simpler alternatives, and the domains differ in which resource binds first. A smartphone keyboard encounters personal vocabulary and changing typing patterns that may be poorly represented in its released model. Local training can use those observations for personalization or contribute updates to a shared model; these are separate objectives. Google's Gboard,[2] for instance, employs federated learning to improve shared models across a large population of users while keeping raw data local to each device (Hard et al. 2018).

The same keyboard also runs two interactive prediction strategies at inference time (figure 11.1). Next-word prediction ranks discrete candidate tokens ("favorite," "address," "schedule") in the suggestion bar, whereas Smart Compose generates multi-word inline ghost text ("ail address") through on-the-fly rescoring directly inside the input buffer. Both mechanisms personalize the immediate interaction through local context; longer-lived model improvement occurs separately when federated training aggregates eligible on-device updates.

Wearable health monitors, which track activity and physiology from accelerometers, heart rate sensors, and electrodermal activity monitors, face a different personalization problem. Individual physiological baselines and sensor placement can affect predictions, motivating calibration or local adaptation. Evaluation must distinguish useful personalization from learning away a change the application is meant to detect. A monitor that adapts to a user's steadily rising resting heart rate has normalized the very signal it exists to flag. Keeping observations local reduces raw-data transfer; whether the adapted model is more accurate is a separate measurement against held-out outcomes.

Voice interfaces in smart speakers and earbuds must recognize commands quickly and accurately in noisy, changing acoustic environments, starting with wake-word detection.[3] These systems face foreground latency requirements, such as an end-to-end response target under 500 ms and a wake-word detection target below 100 ms. Those targets constrain inference and the interference permitted from background learning. Local adaptation addresses a separate quality objective: reducing false activations and missed detections for a user's voice and acoustic environment. Far-field microphone configurations and room acoustics motivate that adaptation, but its value must be measured while preserving foreground response times.

Industrial IoT and remote monitoring systems expose the availability side of the decision. In agricultural sensing, pipeline monitoring, or environmental surveillance, connectivity to centralized infrastructure may be limited, expensive, or unavailable. Local inference permits continued detection while disconnected. Local learning adds the ability to adjust thresholds or track seasonal trends,

2 **Gboard (2017)**: One of the first widely reported commercial federated learning deployments at mobile scale. Published reports showed that keyboard models could improve from aggregated on-device updates while raw typed text remained local. The systems consequence is durable. Compressed updates and privacy mechanisms make cross-fleet learning possible over constrained mobile networks without centralizing raw keystrokes.

3 **Wake-word detection**: An always-on detector incurs inference cost throughout the listening period, even when no keyword occurs. Published keyword-spotting results on dedicated neural processors, plotted in figure 11.6, spend on the order of 30 to 160 μJ per evaluation, so the energy budget is set by how often the detector evaluates over the listening period as much as by the cost of any one evaluation. Local personalization adds a separate update cost and earns its place only if evaluation shows an acceptable trade-off between missed keywords, false activations, and energy use.

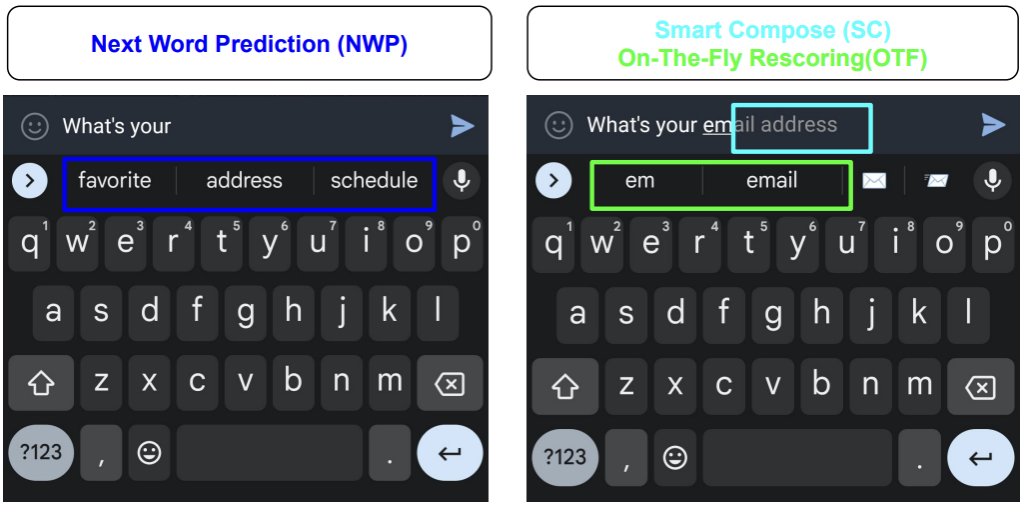

Figure 11.1: On-Device Prediction Strategies: Gboard performs next-word prediction and Smart Compose rescoring locally at inference time. These request-time mechanisms use local context without sending each keystroke to a central server; asynchronous federated training is the separate process that improves shared model parameters from eligible on-device updates. Source: (Xu et al. 2023).

provided the update preserves the distinction between expected variation and anomalies. The consequences of missed detections make validation and a recoverable baseline part of the adaptation design.

Embedded computer vision systems, including robotics, AR/VR, and smart cameras, combine substantial visual processing with the tight foreground deadlines of section 11.2.1, from a 33 ms frame interval down to a sub-10 ms control loop. On-device adaptation can recalibrate a model to new lighting conditions, object appearances, or motion patterns, but it spends resources needed by that foreground pipeline. A useful update must improve validated behavior without exhausting the latency margin that made local inference viable.

Each domain turns a possible adaptation benefit into a resource decision. Mobile keyboards must fit update state alongside typing, wearables must budget energy across repeated updates, voice and vision systems must preserve foreground deadlines, and disconnected sensors must decide whether a fixed model remains adequate.

11.2.3 What crosses the device boundary

Once a domain justifies local learning, the next question is whether each device learns in isolation or shares what it discovers with the fleet. Keeping all learning strictly inside each device prevents a fleet from sharing new words, road hazards, or acoustic patterns discovered across millions of users, yet uploading raw observations to a central cluster exposes private data and consumes wide-area bandwidth. Centralized training collects data on managed infrastructure and releases validated model versions to devices; local adaptation instead changes a device's model using observations that may be unavailable centrally. Hardware, software, power, and connectivity then vary across the learning population, so an operator must coordinate model versions and evaluate updates from intermittently participating clients (Bonawitz et al. 2019).

A returning device may hold an old model or fail to apply a release. Signed updates, functional checks, deployment telemetry, and rollback therefore form part of the learning protocol (Bonawitz et al. 2019).

The central decision is how much state to share across devices. Centralized training, local-only adaptation, and federated coordination are architectural choices that can coexist in one deployment. Local-only adaptation specializes a model to its user's observations, which are often non-IID[4] and noisy, so sharing that specialization requires an explicit boundary between personal state and the parameters intended to improve the population model. Federated coordination adds population

[4] **Non-IID (not independent and identically distributed)**: Samples may be dependent or drawn from different distributions; in a device fleet, user-specific label and feature distributions can make local updates conflict. Such distributions can degrade federated averaging relative to independent and identically distributed (i.i.d.) baselines (Zhao et al. 2018). The observed failure determines whether to change local work, aggregation, or personalization rather than prescribing one remedy for every non-IID population.

learning by periodically aggregating updates to shared parameters. A deployment must specify which parameters enter the aggregate and which remain personal; averaging every local adaptation would erase that distinction. Data locality limits raw-data transfer, while separate privacy controls govern the updates that leave each device. The three architectures therefore differ in what crosses the device boundary, not in how advanced they are (figure 11.2): local-only learning sends nothing, centralized training sends raw examples, and federated learning sends only model updates $\Delta\theta$ while receiving the shared weights θ.

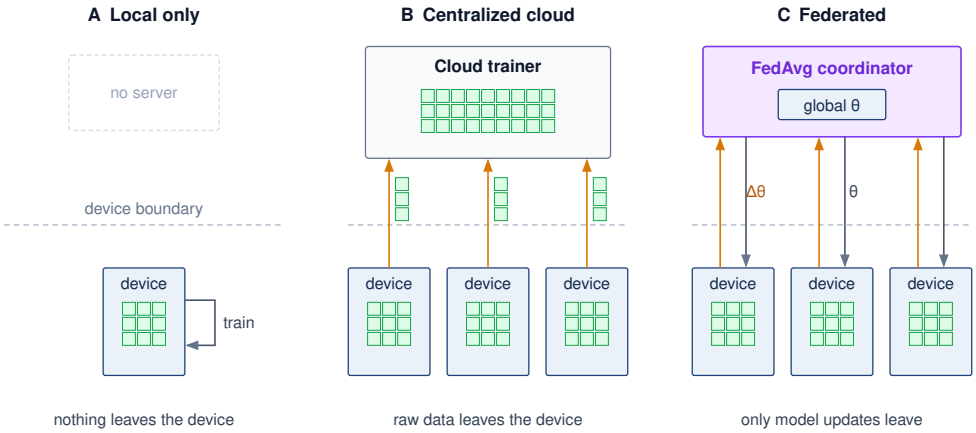

Figure 11.2: Local-Only, Centralized-Cloud, and Federated Learning: The panels compare three architectures by what crosses the device boundary (dashed line), rather than imply a required evolution. (A) Local-only learning trains on the device and sends nothing. (B) Centralized-cloud learning uploads raw data to a cloud trainer. (C) Federated learning keeps data on the device, downloads the global model θ, and uploads only model updates $\Delta\theta$ to a federated averaging (FedAvg) coordinator. Data locality in (C) still requires separate privacy controls for updates and metadata.

The choice of shared and personal state must fit the devices that execute it. Managed data centers can provision capacity and schedule training across dedicated resources; edge devices share limited power, memory, and cooling with foreground applications. Those local budgets determine which adaptation and coordination choices remain feasible.

11.3 Design Constraints

Fine-tuning a neural network on a smartphone is constrained by power and heat: a GPU training cluster may consume megawatts in a controlled facility, whereas a smartphone sustains only a few watts before throttling. Across the edge spectrum, memory ranges from a few gigabytes on phones down to hundreds of kilobytes on microcontrollers, sustained power ranges from a few watts down to milliwatts, and network connectivity may be intermittent or entirely absent. On-device learning must therefore use compute, memory, and data sparingly, and when a local model operates in a safety-critical loop, these resource limits become life-safety engineering problems.

Lighthouse 11.1: Archetype C (Federated MobileNet): The edge power wall

These constraints run through Archetype C (Federated MobileNet), where learning happens on phone-class devices whose power budget spans roughly 2–10 W, a few watts sustained with brief bursts toward 10 W, and which must share that budget with the user's foreground work. Privacy and policy constraints often forbid moving raw examples off the device, and even where transfer is permitted, a mobile uplink and battery rarely afford shipping raw data to the cloud. Data locality is therefore both a policy requirement and a bandwidth and energy constraint. Quantized training, sparse updates, and federated coordination are the strategies that let Archetype C learn within that envelope. Held at that burst ceiling, an idealized 15 Wh battery would be exhausted in about 1.5 hours, and thermal throttling would end the workload

sooner. Microcontroller-class sensors with hundreds of kilobytes of memory and milliwatt budgets sit at an extreme below Archetype C, and this chapter treats them as such.

The constraints on parameters, operations, and data interact. A larger batch can improve hardware utilization while exceeding memory capacity; a smaller trainable state can fit memory while extracting too little signal from local data. Compression provides one way to reach a feasible starting point. Quantization shrinks bytes per weight, pruning removes selected parameters, and knowledge distillation transfers behavior into a smaller architecture. Once the base model fits, local learning must also accommodate the state and work needed to adapt it.

On-device learning adds training state and work to the inference workload. A full-model update requires forward propagation, gradient computation through backpropagation, and weight updates. Memory grows because the update must hold gradients, optimizer state, and the activations saved for the backward pass. Arithmetic grows because the backward pass computes gradients with respect to both activations and weights, up to roughly twice the work of the forward pass. The size of each increase depends on the model, batch, precision, and execution schedule. Compression and selective adaptation address different parts of that pressure, so feasibility requires accounting for the resulting working set rather than applying a universal multiplier.

The relevant budget is the capacity available to the learning task. Centralized training can reserve resources and assemble datasets under a common management policy. A device must fit an update around the operating system, foreground applications, and locally available examples. Total installed memory or peak accelerator throughput therefore overstates what background learning can use.

These constraints define eligibility for a federated round. Model footprint and trainable state determine which device tiers can participate. Sparse, nonuniform data determines how much signal each local update can extract. Limited compute determines when training can run at all. A policy that exceeds the budget on weaker devices changes the population represented in the aggregate, even if it improves throughput on stronger devices.

11.3.1 Quantifying training overhead on edge devices

A compact 10M-parameter model whose 40 MB of FP32 weights fit comfortably inside a smartphone's background-memory allowance can balloon to 360 MB as soon as backpropagation adds gradients, optimizer state, and saved activations. That total exceeds the background memory ceiling and crashes the process on its first backward pass. Training changes memory footprint, memory bandwidth, and hardware utilization through additional live state and backward operations. The illustrative increases in table 11.2 count different things from row to row, so each multiplier carries its own denominator. The 4–12× memory range compares a training working set with the inference weights alone, leaving inference activations and workspace out of the denominator. The 2–3× compute range compares one update (forward plus a backward pass of one to two times the forward work) with one inference. The 10–50× energy range compares a converged adaptation, which takes many gradient steps, with a single inference, so it is not a per-sample multiplier.

Table 11.2: Training Amplifies Inference Constraints: On-device learning operates under the same efficiency constraints as inference but with training-specific amplifications that make optimization dramatically more challenging. This table quantifies how each constraint dimension intensifies when transitioning from running pretrained models to adapting them locally. The factors are illustrative: memory and compute compare an update with inference, while the energy row includes repeated optimization steps toward convergence. These different counting units must not be treated as one universal per-sample multiplier.

Constraint dimension	Inference	Training amplification	Impact on design
Memory Footprint	Model weights + single activation map	Weights + full activation cache + gradients + optimizer state	4–12× the inference weights alone; forces aggressive compression
Compute Operations	Forward pass only	Forward + backward + weight update	2–3× one inference, per update step; limits model complexity
Memory Bandwidth	Sequential weight reads	Bidirectional data flow for gradients	Additional traffic depends on update scope and activation schedule

Constraint dimension	Inference	Training amplification	Impact on design
Energy per Adaptation	Single inference operation	Multiple gradient steps with convergence	10–50× one inference, per converged adaptation; requires opportunistic scheduling
Data Requirements	Current inputs; no training labels required	Local update examples, often sparse or noisy	Necessitates sample-efficient methods
Hardware Utilization	Optimized for forward passes	Different access patterns for backprop	Inference accelerators may not help training

The memory row decomposes into four arrays (figure 11.3). Panel (a) starts with 40 MB of FP32 weights for a 10M-parameter model and adds 40 MB of gradients and 80 MB of Adaptive Moment Estimation (Adam) state for its two moment arrays. Adding 200 MB of assumed activation state brings the total to 360 MB, a 9× amplification over the weights. The same accounting spans the table's range. Weights, gradients, and Adam's two moments alone are four weight-sized arrays, the 4× floor reached when saved activations are negligible. Activations between five and eight times the weight bytes carry the total from 9× to 12×.

Panel (b) holds the 9× proportions fixed across model scales from 1M to 100M parameters, so its constant ratio is an assumption rather than a finding. What it shows is the absolute cost of that ratio. The gap between weights and training stack grows from tens of megabytes at 1M parameters to gigabytes at 100M, so a model whose weights fit comfortably can still have a training stack that does not. Actual peak memory depends on optimizer choice, batch size, activation checkpointing, and how many layers are updated, and the 4–12× range describes these configurations rather than predicting any particular device.

↪ The bytes per parameter that weights, gradients, and optimizer state occupy in mixed-precision Adam training are developed in section A.6.

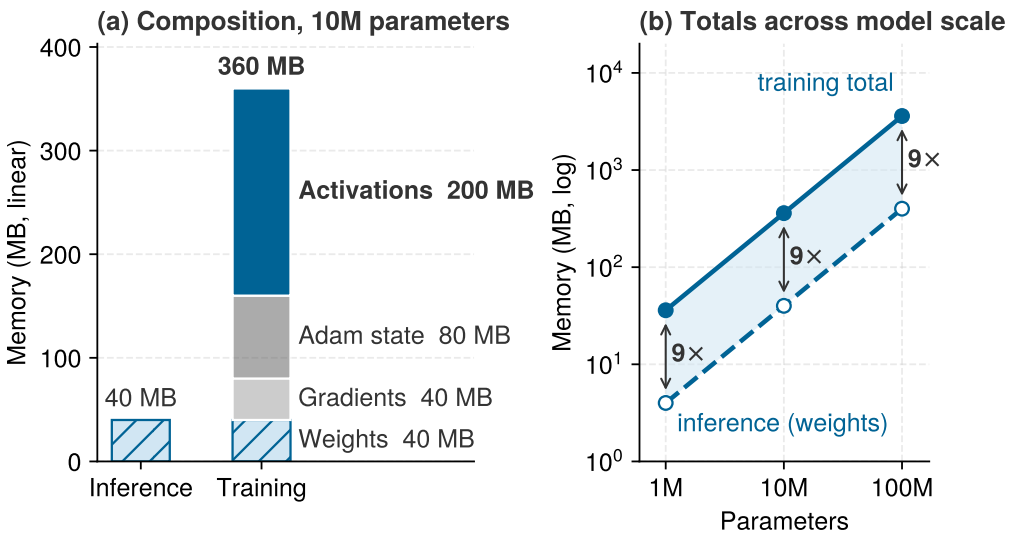

Figure 11.3: The Training Memory Amplifier: (a) On a linear axis, the training stack for a 10M-parameter FP32 model shows what fills memory: gradients match the weights, Adam state doubles them, and activations dominate. (b) On a log axis, inference and training totals keep the same ratio from 1M to 100M parameters. The configuration is illustrative (activations assumed at 5× the weights), producing a representative 9× amplification.

11.3.1.1 Peak memory usage

Consider a smartphone keyboard that uses local gradient updates to adapt to a user's vocabulary and typing patterns. In the illustrative device profile, an update requires 50 MB–100 MB for activations and optimizer state, against a background-memory allowance of 200 MB–300 MB. The actual allowance depends on the operating system, device, and foreground workload. At the selected operating point, the update consumes 25 percent of that allowance before other application state is

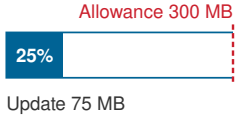

One gradient update claims a quarter of the background-memory allowance.

counted. The design decision is how much learning can fit in the remaining budget while preserving responsive typing; the device's total installed RAM does not answer it.

Whether an update stays inside that allowance depends on peak live residency over the step, because tensors are allocated and released dynamically as execution proceeds. The peak occurs when weights, saved activations, gradients, optimizer state, and temporary workspace coexist, which need not happen at the same operation in every implementation. For the 10M-parameter model of figure 11.3, 40 MB of FP32 weights belong to a 360 MB training working set, of which 200 MB is saved activations. Set against a smartphone's 8 GB of installed RAM, that 360 MB working set looks modest, yet it exceeds even the top of the 200 MB–300 MB allowance that the operating system grants a background task after itself and the foreground applications. Activation checkpointing (also called gradient checkpointing) discards selected intermediate activations and recomputes them during the backward pass (Chen et al. 2016), the same trade that cuts activation memory in data-center training. It trades extra computation for lower activation storage, so its value depends on both the peak it removes and the additional time and energy the device can afford.

↩ Activation checkpointing, which recomputes selected activations in the backward pass to cut activation memory at the cost of extra FLOPs, is established in section 6.3.1.2.

That allowance, not installed RAM, is the ceiling for a working-memory budget that every adaptation strategy in this chapter must meet (figure 11.4). The runtime stages frozen weights from nonvolatile flash into volatile working memory and reserves headroom below the physical device ceiling M_{device} for runtime and application state. Within the working-memory envelope M_{peak}, four quantities contribute: let $M_{\text{weights,resident}}$ be the backbone weights that must reside in working memory, M_{train} the trainable parameters plus their gradients and optimizer state, $M_{\text{activations}}(t)$ the saved activation state, and $M_{\text{workspace}}(t)$ operator workspace. With all terms counted in bytes and disjoint, the required working memory is

$$M_{\text{peak}} = M_{\text{weights,resident}} + M_{\text{train}} + \max_t \left[M_{\text{activations}}(t) + M_{\text{workspace}}(t) \right]. \tag{11.1}$$

Admission then requires $M_{\text{peak}} + M_{\text{reserve}} \leq M_{\text{device}}$, where M_{reserve} holds the runtime, the operating system, and the foreground application; for a background task, $M_{\text{device}} - M_{\text{reserve}}$ is the allowance just described. Weights stored in nonvolatile memory still consume deployment storage, and any staging buffers belong in $M_{\text{workspace}}(t)$. The maximum must follow the actual execution schedule because activation and workspace peaks need not coincide. Each term is a separate lever, and the adaptation strategies of section 11.4 differ in which term they shrink. Freezing weights removes their gradients and optimizer state from M_{train}, but it does not make the frozen base or every backward dependency disappear.

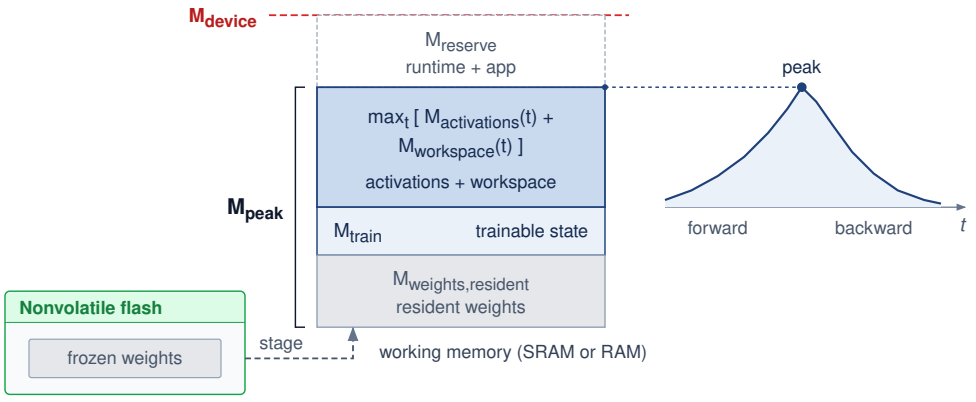

Figure 11.4: Adaptation Memory Budget: Working memory must hold the resident weights $M_{\text{weights,resident}}$, the trainable state M_{train}, and the peak of live activations plus workspace; together these form M_{peak}, which must fit under the device limit once the runtime and foreground application are reserved. Frozen weights kept in nonvolatile flash count toward storage, not working memory, until they are staged, and staging buffers belong to the workspace term. Segment heights are schematic, not to scale.

An inference optimization therefore needs a second feasibility check before it is used for adaptation. Reducing stored weights may leave activations or optimizer state as the dominant cost. Selecting which parameters to train changes that accounting, linking the memory analysis to the adaptation choices that follow.

Checkpoint 11.1: On-device training constraints

Verify understanding of how training amplifies edge constraints:

- ☐ Training memory blowup: Explain where the 4× floor of the chapter's 4–12× training-memory range comes from, and which term pushes a configuration toward the upper end.
- ☐ Timing of the peak: Explain why peak training memory depends on when weights, saved activations, gradients, optimizer state, and workspace coexist, rather than on the sum of every tensor the update ever touches.
- ☐ Backward-pass OOM: Explain why a 10M parameter model can still trigger out-of-memory errors during the backward pass on a smartphone with 8 GB of RAM.
- ☐ Checkpointing energy: Evaluate the claim that activation checkpointing reduces the total energy consumed by on-device training.

11.3.2 Model constraints

A model that fits comfortably on a smartphone can fail before a single backward pass on a smaller edge tier. MobileNetV2 requires approximately 14 MB of storage in its standard configuration. That fits easily in smartphone RAM yet far exceeds the 256 KB of SRAM and 1 MB of flash storage on microcontrollers such as the Arduino Nano 33 BLE Sense.[5] On such severely constrained platforms, even a single convolutional layer may exceed available RAM during training due to intermediate feature maps and gradient storage. Cloud-deployed models can span billions of parameters across multi-gigabyte memory budgets. Models intended for on-device learning must conform to tight limits on memory, storage, and compute, each tightened further by the factors of table 11.2.

Training adds state whose size cannot be inferred from layer count alone. Backpropagation retains or recomputes intermediate activations, with storage determined by tensor shapes, precision, batch size, and the update schedule. The arithmetic grows quickly. A 10-layer convolutional model processing 64×64 images saves feature maps of 64×64 values per channel when its layers keep the input resolution. At FP32, every channel of every retained map then costs 16,384 bytes, about 16 KB. Ten layers of even a few channels each therefore outgrow the 256 KB of microcontroller SRAM before a single gradient or optimizer value is stored. The relevant comparison is that peak against the memory available to learning after runtime and application reserves, which is why the techniques later in this chapter attack saved activations as well as trainable parameters.

Model complexity also affects runtime energy consumption and thermal limits. In smartwatches or battery-powered wearables, sustained training can deplete the energy allowance or trigger thermal throttling even when its tensors fit in memory. Arithmetic precision alone does not determine feasibility; the update duration and measured power determine its energy cost. Benchmarks such as MLPerf Tiny provide small quantized inference targets for constrained devices (Banbury et al. 2021). On-device adaptation techniques address the additional training cost by freezing most parameters, reducing activation storage, or selecting sparse update sets (Cai et al. 2020; Kwon et al. 2024).

Battery and thermal constraints also determine when learning can run. The illustrative smartphone profile uses a sustained processing envelope of 2–3 W and brief bursts of 5–10 W. An update that fits during an idle interval may exceed the available headroom when foreground work starts. Charging supplies energy, but the runtime must still monitor heat and contention, reducing the update duty cycle or suspending learning when limits are approached. An overnight idle-and-charging policy provides one candidate window; measured device conditions determine whether that window remains usable.

These constraints make adaptation feasibility part of architecture selection. When the base model exceeds device capacity, partitioning, quantization, offloading, or a smaller architecture can change the deployment footprint, each with its own execution cost. INT8 or INT4 weights often decide

[5] **Arduino Nano 33 BLE Sense**: With 256 KB SRAM, roughly 65,536× smaller than a flagship smartphone's 16 GB, a single $224\times224\times3$ RGB image occupies ~151 KB and consumes roughly 60 percent of available memory; activation and gradient storage add several more live tensors, and optimizer state can push the training footprint far beyond the inference footprint. A convolutional neural network (CNN) layer can therefore exceed SRAM during backpropagation even when its weights fit. Quantization can reduce this footprint, but feasibility also depends on input size, saved activations, and which parameters are updated.

whether the base model fits at all, and 1-bit weights cut weight bytes 16× relative to FP16 (32× relative to FP32) before scale metadata and activations are counted. Fitting the forward path is the starting point; quantization support and selective-update costs must then be established for the training path.

11.3.3 Data constraints

Fitting the model and its update state into memory solves only the capacity side of the problem, because the available local data still sets how much a fitted update can actually learn. A fitness tracker that records thousands of accelerometer readings across a workout may receive only a single user-supplied activity label, or it may record from a sensor whose calibration has drifted. Unlike the curated datasets used in cloud training, edge data is locally collected, temporally sparse, and often unlabeled. Because sampling frequency, windowing, and temporal correlation can collapse an entire session of correlated sensor readings into a single independent observation, the learning policy must budget for usable supervision rather than raw stored bytes.

On-device data is also frequently non-i.i.d. (Zhao et al. 2018). A voice assistant deployed across households encounters variation in accents, languages, speaking styles, and command patterns. Smartphone keyboards observe individual typing patterns, autocorrect preferences, and multilingual usage. Centralized systems also encounter such variation, but a federated coordinator cannot freely shuffle raw examples across clients. Local updates therefore reflect different data distributions, complicating the balance between shared generalization and personalization.

Label scarcity compounds the distribution problem. Most edge-collected data is unlabeled by default, requiring systems to learn from weak or implicit supervision signals. A smartphone camera may capture thousands of images throughout the day, but only a few are associated with meaningful user actions (tagging, favoriting, or sharing) that could serve as implicit labels. In many applications, including anomaly detection in sensor data and gesture recognition adaptation, explicit labels may be entirely unavailable, making traditional supervised learning infeasible without alternative methods for weak supervision or unsupervised adaptation.

Data quality introduces further challenges. Environmental sensors or automotive ECUs can experience calibration changes, interference, or mechanical wear that corrupt or shift input signals. An adaptation loop may learn from those changes as if they represented useful new behavior. Local input checks and validated outcome signals must distinguish sensor faults from changes the model should learn, especially when raw observations are unavailable for central inspection.

Privacy and security requirements can further restrict which examples may be retained or shared. Keeping health data, personal communications, or behavioral traces on the device narrows raw-data exposure, but local storage and transmitted updates remain sensitive. The learning design must specify access, retention, and permitted disclosures alongside the data needed for adaptation and validation.

11.3.4 Compute constraints

Even when a compact model and validated local examples are in place, the processor that runs inference may not be able to train them: a mobile neural processing unit (NPU) advertised at tens of TOPS for INT8 inference can lack backward-pass or optimizer kernels entirely. An on-device update must then fall back to the CPU and can run far slower and hotter than the chip's inference rating suggests. Across the edge hardware spectrum, from microcontrollers like the STM32F4 and ESP32, through mobile processors with dedicated AI accelerators (Apple Neural Engine, Qualcomm Hexagon, Google Tensor), to high-capability edge modules, inference ratings measure forward-pass throughput on fixed weights. Training workloads add backward and optimizer operations with very different execution paths.

Because an inference rating does not measure this backward path, feasibility depends on tracing where each update operator actually executes. Backpropagation reads saved or recomputed activations, produces gradients, and updates parameters and optimizer state. These operations add memory traffic, with the amount determined by the trainable subset and execution schedule. A specialized accelerator can expand the inference budget without executing any backward operator. The feasibility test is therefore the actual update path. It identifies which processor runs each backward and optimizer operator, including any fallback to the CPU or GPU, and measures time,

memory, and energy for that path under foreground load. Each device class that follows gets this update-path test, applied to its own memory, operators, and power, rather than a verdict read from its inference rating.

At the most constrained end, devices such as the STM32F4[6] or ESP32-S3 microcontrollers offer only a few hundred kilobytes of SRAM and limited compute and power budgets (Warden and Situnayake 2020). Although these families include single-precision floating-point support, practical ML deployments commonly rely on quantized or fixed-point kernels for efficiency. Libraries like CMSIS-NN (Lai et al. 2018) provide optimized neural network kernels for Arm Cortex-M processors, achieving 4.6× runtime improvement through fixed-point arithmetic and single instruction, multiple data optimizations. These severe limitations preclude conventional deep learning libraries and require models designed for quantized arithmetic and minimal runtime memory allocation. Quantization-aware training[7] can prepare models for efficient quantized inference. Executing local updates requires a separately supported training path, potentially with selective updates, quantized gradients, or a compact floating-point working set.

Microcontroller memory rules out training configurations whose live tensors exceed SRAM, rather than every convolutional model. Feasibility depends on tensor shapes, update scope, numerical representation, and the compiled execution schedule. Plain stochastic gradient descent (SGD) reduces optimizer state compared with methods that retain additional moment arrays.[8] Non-neural adaptation, such as k-means clustering, offers a different way to update local centroids for sensor analysis. The method and its arithmetic must still fit the available memory and produce a validated improvement; neither a simple optimizer nor integer execution establishes that result alone.

Moving up the computational hierarchy, mobile-class hardware represents improvement but still operates under severe constraints. Platforms including the Qualcomm Snapdragon, Apple Neural Engine,[9] and the Google Tensor system-on-chip (SoC) provide significantly more compute power than microcontrollers, often featuring dedicated AI accelerators and optimized support for 8-bit or mixed-precision matrix operations. These accelerators offer dedicated matrix multiplication units, on-chip memory hierarchies, and power management features specifically designed for neural network inference, with varying support for local training workloads. While these platforms can support more sophisticated training routines, including full backpropagation over compact models, they still fall far short of the computational throughput and memory bandwidth available in centralized data centers. Training a lightweight transformer on a smartphone, for instance, is technically feasible but must be tightly bounded in both time and energy to avoid degrading the user experience.

Hardware capability contributes to on-device adaptation feasibility alongside the algorithm and runtime. The mobile accelerator scaling tracked in figure 11.5 shows Apple's stated rating climbing from under 1 TOPS on the A11 Bionic to nearly 35 TOPS on the A17 Pro, a step in 2023 that may partly reflect a change in rating precision rather than in the hardware alone. These peak ratings concern forward inference. A training budget must also account for backward operators and the runtime that dispatches them.

Systems Perspective 11.1: The silicon dividend of a dedicated NPU

The gap between running a model on a generic mobile CPU and on a dedicated NPU is best understood as a "silicon dividend" rather than a marginal speedup. Published NPU profiles vary by silicon generation and workload, so this callout uses explicit scenario assumptions rather than a benchmark: a 20× latency improvement and a 60× energy-efficiency improvement relative to the modeled MobileNetV2 baseline (1.883 ms per inference on the CPU).

The dividend belongs to the operators the NPU actually executes. An update adds backward and optimizer work that costs roughly one to two more forward passes, because the backward pass computes gradients with respect to both activations and weights. If the runtime cannot place those operators on the NPU, they fall back to the CPU. The forward pass may still collect its 20× dividend, but CPU backward work and transfers between processors can dominate the complete update. The resulting speed and energy depend on that mixed execution path. An inference-dominated product and an adaptation-dominated one can therefore see very

6 **STM32F4 microcontroller**: With 192 KB SRAM, a 168 MHz clock speed, and a single-precision floating-point unit, the main constraint is not the absence of floating point but the tiny memory and energy envelope. A dense layer with 1,000 input features and 10 hidden units requires approximately 40 KB for weights and biases in FP32, consuming about 20 percent of total SRAM before accounting for activations or gradients. Local updates must therefore fit around the device's primary workload and power budget, and INT8 or fixed-point kernels remain preferable when accuracy allows.

7 **Quantization-aware training (QAT)**: Simulates quantization during training to adapt the model to low-precision inference (Jacob et al. 2018; Krishnamoorthi 2018). Fake-quantization operations commonly retain floating-point training tensors, so QAT does not itself imply integer backward kernels or a smaller training-state budget. On-device adaptation must separately establish supported update arithmetic, state storage, and measured energy costs.

8 **SGD (stochastic gradient descent)**: Plain SGD keeps parameters and gradients, two parameter-sized arrays where Adam keeps four, so its lower state cost must still be weighed against the updates needed to reach the quality target.

9 **Apple Neural Engine**: Apple's dedicated neural accelerator, first shipped in the A11 Bionic in 2017; its rating growth through the A17 Pro is tracked in figure 11.5. Which of an update's operators actually reach it, rather than falling back to the CPU or GPU, is the runtime question worked through in section 11.3.4.3.

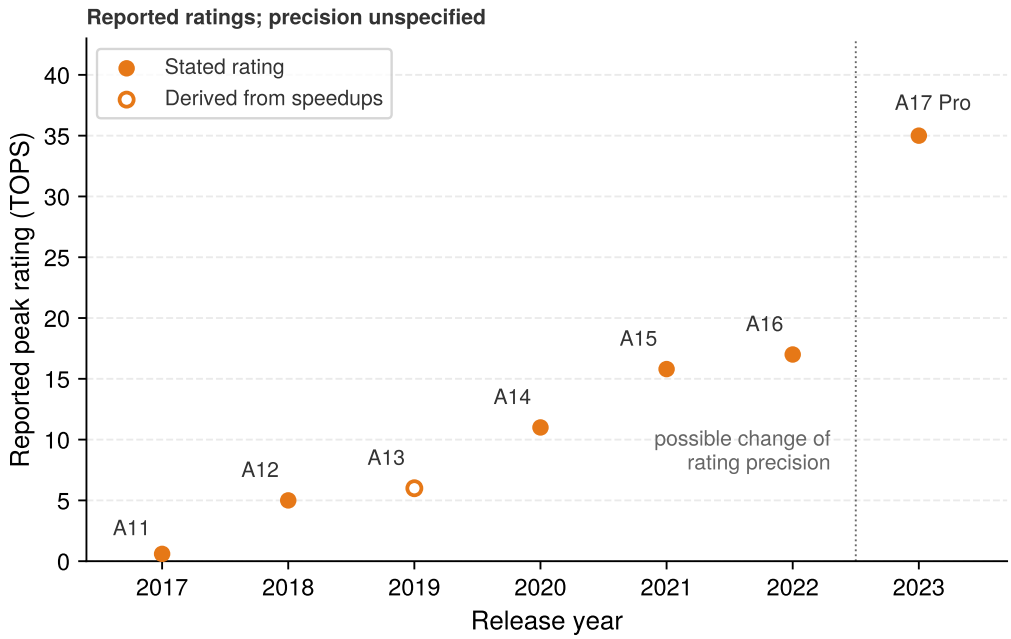

Figure 11.5: Mobile NPU Performance Scaling: Apple's stated peak Neural Engine rating by chip, rising from 0.6 TOPS on the A11 Bionic in 2017 to nearly 35 TOPS on the A17 Pro in 2023. The hollow A13 point is derived from Apple's stated speedups rather than a stated figure. The points are not joined because Apple does not state the rating precision; the 2023 step may partly reflect a change in number format rather than in the hardware alone. These inference-oriented peak ratings do not establish training throughput or backward-operator support. Ratings from Apple (2017), Apple (2018), Apple (2020), Apple (2021a), Apple (2022), and Apple (2023).

different returns from the same silicon, and only the traced update path says which case a deployment is in.

Published benchmark results from MLPerf Tiny and vendor data illustrate the range of edge inference costs in figure 11.6. The plotted tiers span roughly 100× in energy per inference. Dedicated neural processors such as the Syntiant Core 2 and STM32N6 NPU occupy the sub-millijoule MCU and application-specific integrated circuit (ASIC) band. They use 30 to 444 μJ per inference across keyword spotting, visual wake words, and small-image classification, and keyword spotting finishes in under 5 ms. Google Coral appears near 4.8 mJ, and the Jetson AGX Orin appears near 15 mJ for ResNet-50. The roughly 100× gap mixes different workloads with different hardware, so it is not a matched-task efficiency comparison. Energy per inference does not determine battery lifetime on its own, because inference frequency, sleep power, sensing, radio use, and usable battery capacity also enter that budget.

For a learning device the figure sets the scale against which an update must be judged. One evaluation costs microjoules to millijoules, while each second of training inside the 2–3 W sustained envelope spends joules, more than a thousand keyword-spotting evaluations at the plotted 30 to 160 μJ. An update can therefore be a substantial energy cost, but its share of the device budget depends on how often training and inference run. Both paths need measurement over the same operating period.

These computational limitations become especially acute in real-time or battery-operated systems. There, the foreground latency budgets of section 11.2.1 bound how much interference local training may cause, although training itself need not finish within an inference deadline. Training also has different memory access patterns from inference. Gradient computation, activation storage, and optimizer updates add memory traffic whose magnitude depends on which parameters train and which intermediate states are retained. These transfers can create bottlenecks that pure computational

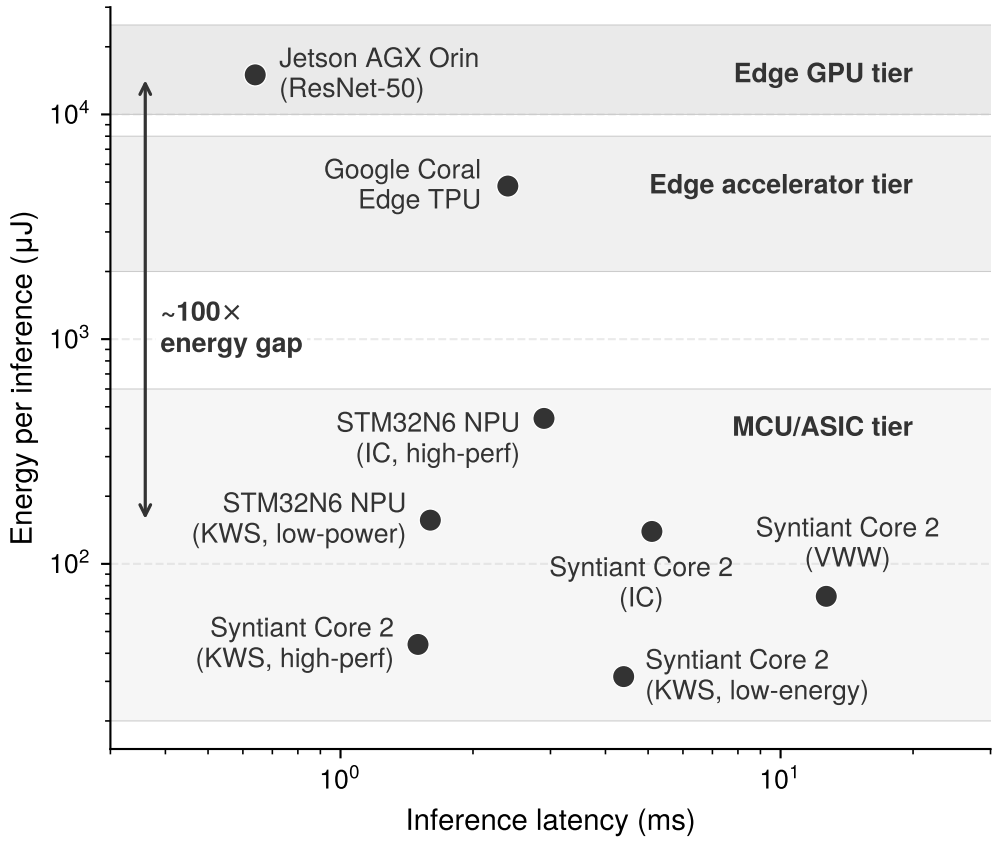

Figure 11.6: The Edge Inference Landscape: Published inference latency and energy measurements from MLPerf Tiny and official benchmarks reveal three distinct hardware tiers separated by approximately 100× in energy per inference; because the tiers run different workloads (keyword spotting and small-image tasks versus ResNet-50), the gap mixes hardware efficiency with workload size. Dedicated neural processors (Syntiant, STM32N6) achieve keyword spotting in under 5 ms at 30–160 microjoules, while edge GPUs (Jetson Orin) deliver sub-millisecond latency at 15 mJ.

metrics do not capture. Hardware designs can address these challenges with specialized support, but a particular device must expose that support through its training runtime. Mixed-precision datapaths may execute different formats in forward and backward kernels when those kernels exist. Sparse-update savings require an execution path that can actually skip the selected work; zero gradients alone do not establish reduced hardware cost (section 11.4.4). Near-memory compute architectures[10] reduce data movement costs by performing gradient updates directly adjacent to weight storage, addressing the memory bandwidth bottleneck. Most edge accelerators, however, remain optimized for inference, which leaves the backward and optimizer operators of local adaptation as an open hardware-software co-design problem.

[10] **Near-memory computing**: Places processing elements adjacent to or within memory arrays to cut the data movement that often dominates arithmetic energy (Horowitz 2014; Kwon and Rhu 2018), which could make write-heavy on-device backpropagation more practical.

11.3.4.1 The mobile memory wall

Rising NPU ratings tempt designers to read device capability from TOPS, but every byte a local update moves crosses the same memory interface as inference traffic, and an update moves far more bytes than an inference. Counting them for the 10M-parameter model of figure 11.3, under the simplest assumption that each tensor crosses DRAM once per use with no cache reuse, shows how far. The forward pass reads the weights and writes the saved activations, and the backward pass rereads the activations and weights and writes the gradients. An Adam step then reads the weights, gradients, and both moments and writes the weights and moments back. One update step moves about 800 MB, 20× the weight bytes a single inference reads. At the 51 GB/s peak bandwidth of an A17 Pro, that traffic alone needs about 16 ms per step before any arithmetic is counted.

Mobile language-model decode, the case already met as the autoregressive bottleneck (principle 18), shows the same ceiling at its starkest, since decode pressure falls on memory and interconnect rather than peak arithmetic throughput (Ma and Patterson 2026). The interface itself is narrow by data-center standards. Data-center HBM3 on an H100 provides 3,350 GB/s of peak unidirectional DRAM memory bandwidth. Flagship mobile systems sit near 51 GB/s (A17 Pro, LPDDR5) to 77 GB/s (Snapdragon 8 Gen 3, LPDDR5X), a gap of roughly 44–65×.

↩ The memory wall, the widening gap between arithmetic throughput and the memory bandwidth that feeds it, is established in section 2.3.

The ratio of bandwidth-bound times equals this bandwidth ratio only when both devices move the same bytes and sustain the same fraction of their peak bandwidth. Precision, batching, cache reuse, and runtime overhead all move the measured ratio. Low-precision weights attack that ceiling when weight traffic is what binds. Storing FP32 weights at 4 bits, for example, reduces weight bytes by 8× and reduces that component of traffic by the same factor if the runtime consumes the compressed representation efficiently. It does not increase physical bandwidth or remove cache-state traffic and decoding overhead, and the precision choice must satisfy both quality and measured latency requirements. For an on-device update the relief is smaller still. In the count above, the three weight reads (forward, backward, and optimizer) are only about 15 percent of the step's traffic, and a quantized base leaves the saved activations, gradients, and optimizer state that make up the rest untouched.

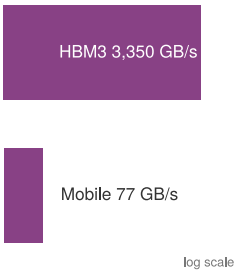

Data-center high-bandwidth memory (HBM) outruns flagship mobile memory bandwidth by about 44 times.

> **Systems Perspective 11.2: Why bandwidth, not TOPS, is the binding constraint**
>
> Small-batch decode is the familiar memory-bound case. Each token rereads weights too large to stay in cache, so the data term $(D_{\text{vol}}/\text{BW})$ of the single-machine iron law sets its time, and sustained bandwidth and bytes per token predict it better than peak TOPS. The roofline recap in section A.3.1 works that decode example. Local adaptation needs the same comparison applied to its own forward, backward, and optimizer traffic, counted per update step rather than per token, where the count above puts it at about 800 MB.

11.3.4.2 Memory hierarchy optimization

Capacity is the question equation 11.1 already answers. Quantization, pruning, and distillation can each shrink the resident and trainable terms of that working set (Jacob et al. 2018; Han et al. 2015; Hinton et al. 2015), but their gains are workload-specific and do not combine automatically, so they must be validated together against the same budget. Once the working set fits, capacity says nothing about the second question, which is how much traffic the update drives through the hierarchy and how much of that traffic the caches can absorb. Operator tiling and memory layout aim to reuse active tensor blocks before eviction, while batch size controls activation storage and available parallelism. Gradient accumulation can limit the microbatch working set, but it also retains accumulation state and adds accesses, so its effect on traffic must be measured.

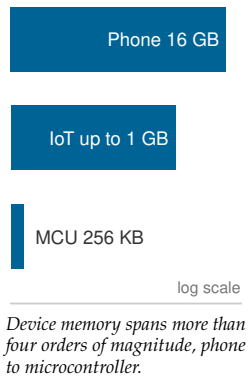

Device memory spans more than four orders of magnitude, phone to microcontroller.

Training adds gradient and optimizer traffic to the tensor reads and writes already required by inference. When these transfers saturate memory bandwidth, additional arithmetic capacity cannot shorten the update. The activation checkpointing introduced with peak memory usage shows why capacity and traffic must be counted separately. Recomputing activations lowers the peak footprint while adding forward work and its memory reads, so capacity can improve as traffic and energy rise. Mixed-precision training changes bytes per stored tensor, but numerical stability and supported backward kernels determine whether the smaller representation is useful. Capacity, traffic, and update time therefore need separate measurements.

11.3.4.3 Mobile AI accelerator optimization

Accelerator choice depends on the training primitives that the device's software stack exposes. Apple Neural Engine, Qualcomm Hexagon DSP, and Google Tensor devices have different compiler and runtime contracts, so the same update can follow different paths on each. A useful execution plan identifies where feature extraction, head or adapter training, and optimizer updates run. A dedicated engine may accelerate the frozen forward path while another processor handles learning, but transfers and duplicate representations then enter the cost. The supported operators and data movement must be established for the actual device and runtime version.

A local update may therefore execute on a CPU or GPU even when inference uses the dedicated neural accelerator. On a real device, the update-path test that opened the discussion of compute constraints becomes a short procedure.

1. List the update's operators in execution order: the forward pass, each backward operator that the trainable subset requires, and the optimizer step.
2. For each operator, record the processor the runtime actually places it on, noting every fallback from the neural engine to the CPU or GPU.
3. At each change of processor, add the transfer and any duplicate tensor representation to the update's memory and traffic accounting.
4. Measure time, peak memory, and energy for the complete traced path while a representative foreground workload runs.
5. Feed those measurements into the peak-memory budget of equation 11.1 and the admission test of equation 11.2 that follows, which together decide whether the update may run.

Mixed-precision datapaths, sparse-update support, and near-memory compute can reduce those costs, but none of them is a prerequisite for training. What matters is whether the available runtime can execute the required operators within the device envelope.

11.3.5 Energy, heat, and background admission

An update path that fits in memory and runs on supported kernels still fails if 30 minutes of background backpropagation noticeably drains the battery, warms the handset in the user's pocket, or delays foreground touch responses. Because energy draw, thermal dissipation, and shared-processor contention reach the user directly, a mobile power budget constrains both the model complexity a device can train and the schedule on which background learning is admitted.

> Napkin Math 11.1: Battery drain: The cost of edge learning
>
> **Problem**: A team is designing a background fine-tuning job for a personalized voice assistant on a smartphone. Run unconstrained, the training job draws 4.5 W, above the 2–3 W sustained envelope given earlier, and takes 30 minutes to complete. If the phone has a 15 Wh battery, how much of the user's battery will this "invisible" update consume?
>
> **Math**:
>
> 1. Total energy: 4.5 W $\times$ 0.5 h = 2.25 Wh.
> 2. **Battery impact**: 2.25 Wh/15 Wh = 15 percent.
>
> **Systems insight**: Consuming 15 percent of a user's battery for one background task makes update frequency a deployment decision. That energy cannot be translated into a fixed amount of screen-on time without specifying the foreground workload. A policy of **opportunistic scheduling** admits learning when device conditions permit it; a federated policy might require charging, Wi-Fi, and thermal headroom. Charging supplies energy but does not remove heat or foreground contention, so those constraints still need separate checks.

This battery calculation is the local edge-learning version of a broader energy constraint. Training must satisfy the user's energy budget in addition to the model's accuracy target. A production design therefore has to account for energy at scheduling time, compare local computation against communication cost, and defer learning when the user's device cannot absorb the update.

That decision can be expressed as a background-work admission condition. Let E_{update} be the predicted energy for one adaptation window, E_{slack} the energy the policy permits before the next charging opportunity, P_{device} the current sustained power, P_{update} the adaptation increment, and P_{thermal} the thermally sustainable limit. Let $L_{\text{lat,fg}}$ be measured foreground latency, $\Delta L^{\text{wc}}_{\text{lat,fg}}$ the worst-case latency increase under concurrent adaptation, and T_{deadline} the foreground deadline. The runtime admits an update only if

$$E_{\text{update}} \leq E_{\text{slack}} \quad \wedge \quad P_{\text{device}} + P_{\text{update}} \leq P_{\text{thermal}} \quad \wedge \quad L_{\text{lat,fg}} + \Delta L^{\text{wc}}_{\text{lat,fg}} \leq T_{\text{deadline}}. \tag{11.2}$$

Charging can enlarge the energy allowance, but it does not waive the thermal or foreground-latency tests. The quantities must come from device- and workload-specific profiling, and the runtime must revoke admission when the measured envelope changes. Shorter update windows, a lower duty cycle, or a smaller trainable parameter set can make a rejected update feasible; allowing background learning to consume the foreground deadline cannot.

The same test also decides who teaches a shared model. Passing it correlates with hardware class, charging habits, and connectivity, so devices that are old, rarely charged, or thermally tight contribute few updates. As a result, the population a federated model learns from is the admitted population rather than the deployed one (principle 20). The admission policy that protects each device is therefore also a sampling policy, and section 11.7.1 measures the skew it creates.

Example 11.1: On-device keyboard learning

Scenario: Consider a mobile keyboard that trains its next-word predictor locally and contributes updates to a shared model through federated averaging (McMahan et al. 2017).

Diagnosis: In the keyboard profile of section 11.3.1.1, one local update needs 50 MB–100 MB of working memory against a background allowance of 200 MB–300 MB. At the selected operating point, the update alone therefore claims 25 percent of the allowance before the keyboard's own state is counted, and the resident weights in M_{peak} claim more. Exceeding the allowance risks process eviction mid-update. Energy is the second constraint. An unconstrained job like the one in notebook 11.1 would take 15 percent of the battery, more than a policy that protects the user's charge is likely to grant as E_{slack} in equation 11.2 while the phone runs on battery. The pipeline therefore admits training only when the device is idle, charging, and on Wi-Fi, and it quantizes update vectors to shrink the transmitted payload D_{vol}.

Systems lesson: A federated keyboard must satisfy all three terms of equation 11.2 (energy $E_{\text{update}} \leq E_{\text{slack}}$, sustained power under P_{thermal}, and typing latency under its deadline), fit M_{peak} within its memory allowance, respect uplink bandwidth BW, and meet its declared privacy guarantees together. Differential privacy adds a specific release guarantee and its associated costs; it is not implied by keeping raw text local.

11.3.6 Matching edge constraints to adaptation pillars

Memory, data, compute, and admission constraints act on the same update window at once. Shrinking trainable parameters so an update fits in RAM still fails if local supervision is too sparse to guide the step or if intermittent connectivity drops the device from every federated round. The constraint analysis therefore yields three linked design decisions. The training amplification of table 11.2 motivates measuring the update footprint against equation 11.2 and choosing which parameters to adapt. Information scarcity, including limited local datasets, non-IID distributions, restrictions on data sharing, and minimal supervision, motivates data efficiency methods that reuse available evidence. When shared learning is useful, device heterogeneity and intermittent availability motivate federated coordination with explicit validation and privacy controls. Each pillar answers a different constraint within the same deployment budget and brings its own techniques to that answer (table 11.3).

Table 11.3: Constraint-Solution Mapping: The three fundamental constraint categories in on-device learning each drive corresponding solution approaches through direct necessity.

Constraint category	Key challenges	Solution approach
Resource Amplification	• Training workloads (4–12× memory) • Memory limitations • Power constraints	**Model Adaptation** • Parameter-efficient updates • Selective layer fine-tuning • Low-rank adaptations
Information Scarcity	• Limited local datasets • Non-IID distributions • Privacy restrictions	**Data Efficiency** • Few-shot and streaming adaptation • Experience replay • Compressed representations

Constraint category	Key challenges	Solution approach
Coordination Challenges	• Device heterogeneity • Intermittent connectivity • Distributed validation complexity • Scalability requirements	**Federated Coordination** • Federated averaging and asynchronous aggregation • Secure aggregation and differential privacy • Client selection and straggler handling

11.4 Model Adaptation

The 10M-parameter model of figure 11.3 is an inference candidate for a smartwatch with 500 MB of total system memory, because its FP32 weights occupy 40 MB. Training every parameter would need the full 360 MB stack, most of the watch's memory before its operating system and foreground application take their share. The weight footprint alone does not establish inference fit; activation buffers, runtime workspace, and foreground reserve must also fit. Full-model training leaves much less room for that remaining state. The system therefore trains only a compact, strategically placed parameter set and, where it can, also shrinks the activations that the backward pass must keep. Resource-efficient model adaptation then determines both the local work a device can complete and the shared update format the fleet can aggregate.

The engineering choice balances adaptation expressivity against resource consumption. Full fine-tuning exposes all parameters to learning and can fit sufficiently small models; retaining a fixed model can be adequate when local context already supplies the needed personalization. Between these choices, selective updates offer smaller training state. The device resource envelope determines which candidates fit: a wearable-class device with 1 MB RAM has a different budget from a smartphone with 8 GB. The local variation then determines how much trainable capacity is useful. Bias updates may capture calibration changes, while shifts requiring new features need a richer update space. Runtime support, inference cost, shared-update compatibility, and rollback determine whether the candidate can enter the deployed system.

The selection of adaptation techniques follows this matching logic, from weight freezing (section 11.4.2) to structured and sparse updates (section 11.4.4). Their order of presentation is not a universal ranking of cost or expressivity; placement, update size, and runtime support determine those properties. Building on the resource accounting in section 11.3, the target is the smallest update space that captures the local variation: *preserve what works globally, adapt what matters locally*. Weight freezing restricts training to selected existing parameters, such as bias terms or final layers, whereas structured updates add low-rank or residual components to a frozen base and sparse updates select existing parameters according to their adaptation value and measured cost. Each choice must preserve enough useful capacity while fitting the device's update budget.

11.4.1 Transfer as the premise

All three adaptation families work only because the device does not learn its representation from scratch. Training a feature extractor from random initialization on a few dozen noisy phone interactions would overfit immediately and exhaust the energy allowance before learning anything generalizable. Instead, every local update that fits an edge budget relies on knowledge transfer, in which a cloud-pretrained model supplies representations useful for a new task or domain. Reusing those representations reduces what the device must learn from local examples, and the benefit depends on what the tasks share. A model pretrained in the cloud can adapt with limited local updates when its representations retain the distinctions the new task requires. If a new input modality or objective requires features the base model does not provide, a small local update may be insufficient.

Transfer is therefore the premise of the adaptation methods that follow. When pretrained representations carry the distinctions a personal typing vocabulary, an individual gait, or a drifting sensor needs, bias-only updates, adapters, and head-only training (which retrains only the final classification layer) can learn a small correction while freezing most of the network. When that premise fails, head-only adaptation cannot recover distinctions absent from its fixed features. Earlier-layer adapters or broader updates may learn them, but their added state, computation, and data requirements must fit the device budget.

11.4.2 Weight freezing

When pretrained features already capture the structure of the task, the most direct way to shrink M_{train} in equation 11.1 is to freeze the majority of a model's parameters and adapt only a minimal subset. **Bias-only adaptation** holds all weights fixed and updates only the bias terms (scalar offsets applied after linear or convolutional layers). The constraint reduces trainable parameters by 100–1000×, removes the weight gradients and their optimizer state from M_{train}, and helps mitigate overfitting when training data is sparse or noisy.

Consider a standard neural network layer:

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

where $\mathbf{W} \in \mathbb{R}^{m \times n}$ is the weight matrix, $\mathbf{b} \in \mathbb{R}^m$ is the bias vector, and $\mathbf{x} \in \mathbb{R}^n$ is the input. In full training, gradients are computed and applied to both $\mathbf{W}$ and $\mathbf{b}$. In bias-only adaptation, the system freezes the weight tensor ($\Delta\mathbf{W} = \mathbf{0}$) and computes gradients exclusively for the bias:

$$\mathbf{W} \leftarrow \mathbf{W}, \quad \mathbf{b} \leftarrow \mathbf{b} - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{b}}$$

This reduces the number of stored gradients and optimizer states, enabling training to proceed under memory-constrained conditions. The smaller trainable state helps only if the device also supports the arithmetic and activation schedule required by the update. In the bias-only setup of listing 11.1, all convolutional and fully connected weights remain frozen, and only the bias terms adapt to local data.

Listing 11.1: Bias-Only Adaptation: Freezes model parameters except for biases to reduce memory usage and allow on-device learning.

```
# Freeze all parameters
for name, param in model.named_parameters():
    param.requires_grad = False

# Enable gradients for bias parameters only
for name, param in model.named_parameters():
    if "bias" in name:
        param.requires_grad = True
```

This pattern restricts optimizer updates to bias terms. What it saves, and what it does not, follows from the two derivatives of the layer above. The weight gradient $\partial\mathcal{L}/\partial\mathbf{W} = (\partial\mathcal{L}/\partial\mathbf{y})\,\mathbf{x}^T$ needs the layer's saved input $\mathbf{x}$, whereas the bias gradient $\partial\mathcal{L}/\partial\mathbf{b} = \partial\mathcal{L}/\partial\mathbf{y}$ needs only the gradient arriving from above. A frozen linear layer therefore no longer has to keep its input for its own sake. Reaching a bias in an earlier layer, however, still sends the gradient back through every later layer. Each nonlinearity or normalization on that path needs its own saved state (a ReLU needs at least the mask of which inputs were positive). Freezing parameters thus shrinks M_{train} without eliminating backward computation or all saved activation state. This is valuable when adapting pretrained models to user-specific or device-local data where the core representations remain relevant but require calibration.

Tiny Transfer Learning (TinyTL) (Cai et al. 2020) targets the activation term of equation 11.1 as well as trainable state. It builds on the derivative argument above. Because frozen linear and convolutional layers drop their saved inputs, the remaining activation footprint comes from nonlinearities, normalization layers, and any trainable side-paths added to recover expressivity. Adding a standard residual adapter beside a frozen block would require saving that block's full-resolution input feature map to compute the adapter's weight gradients, undoing the memory saving. TinyTL avoids that trap with lite residual modules that downsample the input feature map and project it to a narrow channel width before applying trainable convolutions, so the backward pass caches only a small low-resolution tensor at each block alongside compact group-normalization and activation state.

Full fine-tuning (Panel A of figure 11.7) computes gradients for every intermediate layer weight and caches every block's full input map, whereas TinyTL (Panel B) freezes the convolutional backbone

(gray blocks) and confines parameter updates to the downsampled lite residual side-paths and bias vectors **b** (blue blocks). The dashed backward-pass arrow spans the full network depth in both architectures. TinyTL's memory reduction therefore stems not from truncating the backward pass, but from replacing full-resolution saved feature maps with the much smaller inputs of the lite residual modules.

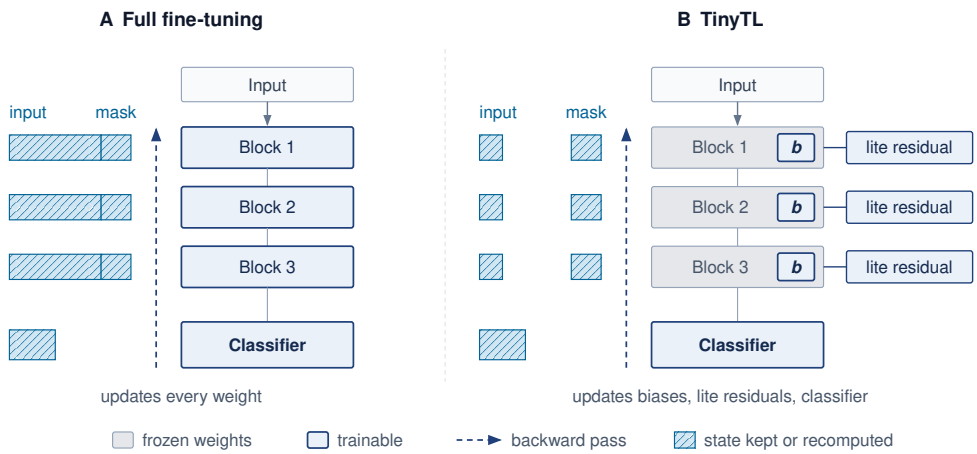

Figure 11.7: TinyTL Adaptation Paths: (A) Full fine-tuning updates every weight and retains the state needed for their gradients. (B) TinyTL freezes the backbone weights (gray) and trains only each block's bias vector (**b**), small lite residual modules beside the blocks, and the classifier (blue). The backward pass still traverses every frozen block to reach those biases, so the saving comes from what must stay live for it, not from a shorter backward path; the activation and normalization choices determine that state.

The memory reduction is useful only when the frozen backbone remains expressive enough for the downstream task. The bias terms allow for minor but meaningful shifts in model behavior, particularly for personalization tasks. When domain shift is more significant, TinyTL can optionally incorporate small residual adapters to improve expressivity, all while preserving the system's tight memory and energy profile.

TinyTL demonstrates that reducing saved activation state can matter more than reducing the number of trainable parameters alone (Cai et al. 2020). On its image-classification transfer benchmarks, TinyTL reports training-memory savings of up to 6.5× relative to fine-tuning the full network with little accuracy loss, so the adaptation working set can shrink to roughly 15 percent of the full fine-tuning footprint. The lite residual modules that recover adaptation capacity add only about 3.8 percent memory overhead in that study. Those savings depend on the selected feature extractor, input resolution, batch size, normalization, and residual-module configuration. A saving measured on those benchmarks says nothing about whether an arbitrary frozen backbone fits into microcontroller SRAM, and when a deployed model must learn richer feature combinations without custom downsampled side-paths, it needs general structured updates beside the frozen weights.

11.4.3 Structured parameter updates

Adapting a smartphone camera pipeline to a new low-light sensor profile or a mobile language model to a user's technical vocabulary requires recombining features across channels, a change that scalar bias offsets **b** cannot express while the weight matrices **W** remain fixed. Unfreezing **W** directly would restore the full gradient and optimizer arrays in M_{train}, so the system instead adds compact, structured trainable matrices beside the frozen backbone. Residual adaptation modules (Houlsby et al. 2019) and low-rank parameterizations (Hu et al. 2021) implement this choice in different ways. Both restrict parameter updates to the added components, but their training cost still includes the backward path through the surrounding network. The comparison with bias-only or full fine-tuning therefore requires both a quality target and a peak-memory measurement.

11.4.3.1 Adapter-based adaptation

A common implementation involves inserting residual adapters, which are small residual bottleneck layers, between existing layers in a pretrained model. Consider a hidden representation $\mathbf{h}$ passed between layers. A residual adapter introduces a transformation:

$$\mathbf{h}' = \mathbf{h} + A(\mathbf{h})$$

where $A(\cdot)$ is a trainable function, typically composed of two linear layers with a nonlinearity:

$$A(\mathbf{h}) = \mathbf{W}_2\,\sigma(\mathbf{W}_1\mathbf{h})$$

with $\mathbf{W}_1 \in \mathbb{R}^{r\times d}$ and $\mathbf{W}_2 \in \mathbb{R}^{d\times r}$, where $r \ll d$. This bottleneck design ensures that only a small number of parameters are introduced per layer.

The adapters act as learnable perturbations on top of a frozen backbone. Because they are small and sparsely applied, their weights add little storage, although their saved activations and optimizer state still count against the training budget. They let the model shift its predictions in response to new inputs.

The residual bottleneck in listing 11.2 maps onto the adapter equation term by term. Its down-projection plays the role of $\mathbf{W}_1$, mapping width d to the bottleneck width r. A ReLU supplies σ, and the up-projection plays the role of $\mathbf{W}_2$, restoring width d before the input is added back. Such modules can be implemented efficiently on edge devices, particularly when the two projection matrices are small and fixed-point representations are supported.

Listing 11.2: Residual Bottleneck Adapter: The code implements a compact adapter module with down-projection and up-projection layers, reducing trainable parameter count while enabling efficient model adaptation on edge devices.

```
import torch.nn as nn


class Adapter(nn.Module):
    def __init__(self, dim, bottleneck_dim):
        super().__init__()
        # Project from full dimension to bottleneck (e.g., 768 -> 16)
        self.down = nn.Linear(
            dim, bottleneck_dim
        )  # W1: learns compression
        # Project back to original dimension (e.g., 16 -> 768)
        self.up = nn.Linear(
            bottleneck_dim, dim
        )  # W2: learns expansion
        self.activation = nn.ReLU()

    def forward(self, x):
        # Residual connection: original + nonlinear bottleneck
        # Only adapter params trained; base model frozen
        return x + self.up(self.activation(self.down(x)))
```

11.4.3.2 Low-rank techniques

Another efficient strategy is to constrain the weight update itself to low rank. Low-rank adaptation (LoRA)[11] leaves a base matrix $\mathbf{W}_0 \in \mathbb{R}^{m\times n}$ frozen and learns its change as the product of a down-projection $\mathbf{A} \in \mathbb{R}^{r\times n}$ and an up-projection $\mathbf{B} \in \mathbb{R}^{m\times r}$, with $r \ll \min(m,n)$. The adapted layer adds $\mathbf{BAx}$ to the frozen output $\mathbf{W}_0\mathbf{x}$. Training cost follows from that structure. The update $\Delta\mathbf{W} = \mathbf{BA}$ has $r(m+n)$ trainable parameters instead of mn, and only those parameters carry gradients and optimizer state.

The decomposition rests on a linear-algebra fact. Any matrix can be expressed as a sum of rank-one matrices through singular value decomposition. Constraining updates to low rank (typically $r = 4$ to 16) captures a restricted set of variation modes while reducing parameters. For a typical transformer layer with dimensions 768×768, full fine-tuning requires updating 589,824 parameters. With rank-4 decomposition, only $768\times4\times2 = 6{,}144$ parameters are updated, a 99 percent reduction. Whether

[11] **LoRA (low-rank adaptation)**: Hu et al. (2021) learn rank-r decomposition matrices $\mathbf{A}$ and $\mathbf{B}$ instead of updating full weight matrices, reporting a 10,000× reduction in trainable parameters for their GPT-3 comparison. That reduction does not apply to the full model footprint: a model with 7B parameters has about a 14 GB FP16 weight footprint before gradients, optimizer state, and activations. Even at rank 16, adapter storage depends on the dimensions and number of adapted matrices, and deployment must still accommodate the frozen base model.

↩ The low-rank adapter forward pass, in which a frozen base projection is summed with a trainable down- and up-projection, is established in section 10.7.9.1.

this retains most of the adaptation quality depends on the task, rank, and base model; the LoRA authors report that low-rank adapters can match or approach full fine-tuning on several transformer adaptation benchmarks (Hu et al. 2021).

LoRA was developed for transformer models but applies to any layer built on a weight matrix. From a systems engineering perspective, LoRA makes the trainable adapter the artifact that can be stored, transmitted, scheduled, and rolled back separately from the frozen base model.

Drawn to proportion for a square layer of width d, the adapted layer shows where the saving comes from (figure 11.8). Input $\mathbf{x} \in \mathbb{R}^d$ passes through the frozen $d \times d$ base matrix $\mathbf{W}_0$ and, in parallel, through a thin $r \times d$ down-projection $\mathbf{A}$ and a $d \times r$ up-projection $\mathbf{B}$. The two paths sum to $\mathbf{h} = \mathbf{W}_0\mathbf{x} + \mathbf{BAx}$. The base weights need no parameter gradients or optimizer state, although backward computation may still traverse the base path. Folding the product into the base weights ($\mathbf{W}_0 + \mathbf{BA}$) removes separate adapter arithmetic for that merged model. Numerical compatibility, any requantization, and the cost of retaining merged copies determine whether this is the appropriate serving representation.

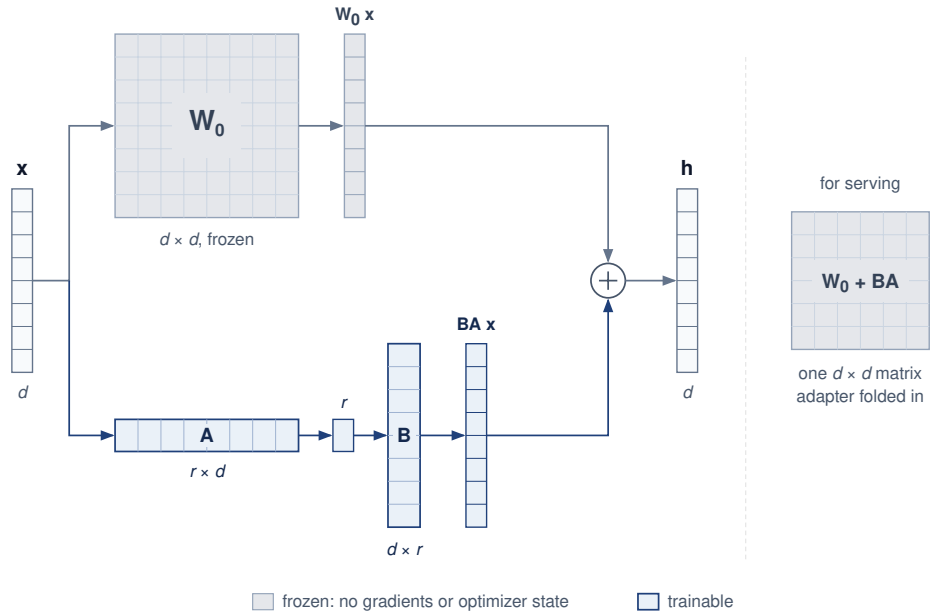

Figure 11.8: LoRA Architecture: The input $\mathbf{x}$ passes through the frozen base matrix $\mathbf{W}_0$ and, in parallel, through a trainable low-rank path: down-projection $\mathbf{A}$ ($r \times d$) and up-projection $\mathbf{B}$ ($d \times r$), whose product is the low-rank update $\Delta\mathbf{W}$. Only $\mathbf{A}$ and $\mathbf{B}$ carry gradients and optimizer state. For serving, the product $\mathbf{BA}$ can be folded into $\mathbf{W}_0$, removing the separate adapter operations when the merged representation is supported. End-to-end latency still needs measurement.

Consider a mobile deployment where a language model with 7B parameters has a 14 GB FP16 weight footprint, before accounting for the additional gradients, optimizer states, and activations required by full fine-tuning. That footprint exceeds smartphones with about 8 GB total memory unless the base model is quantized, partitioned, or offloaded. Choosing LoRA rank 16 restricts the update space, and this chapter's fleet examples budget a reference adapter update of 50 MB. That figure is a communication payload, not a training-memory estimate. The device must also fit the base model, gradients, optimizer state, and live activations under its execution plan.

The smaller payload also matters under intermittent connectivity. Transferring the full FP16 model in this example requires 14 GB, while the reference adapter update is 50 MB. Transfer time depends on encoded bytes, effective link throughput, and interruptions. An adapter-only transfer provides the saving only when the receiving device already holds a compatible base model; base-version changes must be included in the communication budget. In both the residual bottleneck and LoRA, the width or rank r sets the trainable cost, so a deployment can size it to the memory each device

class leaves for learning. The device-tier framework of section 11.4.3.4 decides which mechanism each class receives.

11.4.3.3 Edge personalization

Adapters support multiple personalized contexts with one frozen backbone. For an input width d, a residual bottleneck of width r uses $2dr$ matrix parameters, plus $r + d$ biases when both projections include them. A deployment with several adapters must sum this count over their insertion sites; rank or bottleneck width alone does not determine the total. Stored bytes then equal parameter count times bytes per stored value. With $d = 768$ and $r = 16$, one adapter holds 25,360 parameters, about 50.7 KB in FP16; one adapter after each of 12 blocks totals 304,320 parameters, roughly 0.6 MB. During training, gradients, optimizer state, and saved activations add to that storage footprint. A device can swap context-specific modules while sharing the much larger base model, provided it can hold the base and active adaptation state together.

This **adapter switching** pattern lets a smartphone select context-specific behavior without replacing the entire backbone. When the user enters a low-light environment, the system can load a compatible adapter into the vision pipeline; selecting an existing adapter does not itself require learning. The modularity extends naturally to federated learning: transmitting only the adapted parameters can reduce upload traffic relative to a full-model update, with savings determined by the total adapter count and encoding. Residual-adapter studies show that small trainable modules can adapt a shared visual backbone across multiple visual domains (Rebuffi et al. 2017), making them a useful design point between personalization quality and system resources. Freezing the backbone preserves its parameters, while adapters provide a bounded place for user- or context-specific updates. The resulting model can still suffer catastrophic forgetting, the loss of previously learned behavior after new updates (defined in section 11.5.2).

Napkin Math 11.2: The hidden cost of personalization

Problem: A team is deploying a 10M parameter vision model to a smartphone with support for 10 different "User Contexts" (Home, Office, Car, etc.). If a full fine-tuned model requires 40 MB, how much storage does using residual adapters save instead?

Math:

1. **Full fine-tuning**: 10 contexts $\times$ 40 MB = 400 MB.
2. **Adapter approach**: (1 $\times$ 40 MB backbone) + (10 $\times$ 200 KB adapters) $\approx$ 42 MB.
3. **Storage savings**: 400 MB/42 MB $\approx$ 9.5$\times$ total device storage reduction.
4. **Per-context efficiency**: 40 MB/0.2 MB = 200$\times$.

Systems insight: Multiple personalized models create a storage cost even when only one runs at a time. Storing 10 copies of a 40 MB model consumes 400 MB. In this example, sharing the frozen backbone and storing separate adapters reduces the marginal storage cost of a new context by 200$\times$. This comparison establishes storage savings; the device must still accommodate the active model, adaptation state, and context-switching cost.

In smartphone camera pipelines, environmental lighting, user preferences, and lens distortion vary between users. Residual modules can adapt a frozen shared model to these differences while leaving its base weights unchanged; the combined model's behavior still requires validation. In voice-based systems, adapters similarly provide a way to test personalized speech recognition without retraining the full acoustic model. Keeping the base and adapter versions compatible also supports rollback or switching between user-specific versions when adaptation degrades performance.

11.4.3.4 Performance vs. resource trade-offs

Selecting an adaptation strategy requires accounting for both resource cost and expressivity. Bias-only updates retain a small trainable state but may lack the capacity to handle structural domain shifts. Full fine-tuning exposes every parameter to adaptation and adds corresponding gradient and optimizer state. LoRA and residual adapters offer intermediate choices whose cost depends

on rank or width, insertion sites, and numerical representation. The same budget, equation 11.1, governs each of these strategies. Because peak working memory M_{peak} must jointly contain resident weights, trainable state M_{train}, and the peak of live activations under the device ceiling M_{device}, each strategy shrinks a different term. Bias-only and adapter methods shrink M_{train}, and TinyTL-style choices also shrink the activation term. Training only the final classifier, whose backward pass stops at the head, removes most of that term. Parameter storage alone therefore cannot determine whether background adaptation fits.

LoRA and residual adapters impose different deployment costs. For a fixed adapter and compatible numerical representation, LoRA's linear update can be merged into the backbone before inference ($\mathbf{W}_0 + \mathbf{BA}$), avoiding separate adapter arithmetic in the forward path. A quantized base may require requantization and renewed quality validation after merging. Keeping a separate merged backbone for every context also sacrifices the storage benefit of sharing one base. Residual adapters remain distinct modules ($\mathbf{y} = f(\mathbf{x}) + \text{Adapter}(\mathbf{x})$); their latency cost depends on shape, placement, and runtime, while switching contexts requires only the corresponding adapter state.

The runtime must support the selected adapter's forward operators and its training path. These can use a dynamic graph or a statically compiled update schedule; dynamic graph construction is not itself a requirement. For an inference engine such as Open Neural Network Exchange (ONNX) Runtime Mobile, a runtime that executes an adapter's forward path may still lack its backward operators, optimizer, or state management. The deployment must verify both paths, including any transfers between the processor that runs inference and the processor that performs updates.

For this multi-context personalization workload, where one device switches among several adapters, system architects can use a device-tier framework to organize candidate mechanisms. On flagship phones (Tier 1) with dedicated Neural Engines and 8 GB of RAM or more, residual adapters provide a strong balance of modularity and speed, because switching contexts swaps only adapter state. On mid-range devices (Tier 2) with 1 GB up to 8 GB, LoRA may fit when a fixed adapter can be merged into a compatible base representation before inference, trading switching flexibility for zero added inference arithmetic. For ultra-constrained IoT endpoints (Tier 3) with less than 1 GB RAM, bias-only updates may be the only practical option.

Budget phones show that the tier follows the memory available to learning rather than installed RAM. Their 4 GB places them in Tier 2, yet only 1 GB–2 GB remains for ML work after the operating system. A background adaptation task may be granted less, so a deployment can reasonably hold them to the Tier 3 mechanism. The assignments follow from the workload as much as from the hardware. A workload that keeps one adaptation per user, rather than many switchable contexts, loses the reason to keep adapters separate and can favor merged LoRA even on flagship devices. The runtime must still verify the base-model footprint and admit training within the measured device envelope.

Checkpoint 11.2: Edge personalization

Verify understanding of efficient on-device adaptation:

- ☐ Adapter switching: Explain why Adapter Switching is more storage-efficient than maintaining separate full-model copies for different user contexts.
- ☐ Adapter size: Derive the parameter count of a residual adapter from its input width and bottleneck width, then identify the additional inputs needed to compare its storage with a full model.
- ☐ Residual vs. LoRA: For a multi-context workload, identify the deployment tier (Tier 1, 2, or 3) where residual adapters are preferred over LoRA and justify the choice, then explain why a workload with one adaptation per user changes the answer.
- ☐ Adapter bandwidth: Evaluate the claim that transmitting adapter weights in federated learning consumes more bandwidth than transmitting the entire model.

11.4.4 Sparse updates

Bias-only updates and fixed adapters choose where the model may change before seeing local data, yet a local shift can alter specific convolutional filters or attention channels rather than every block uniformly. Sparse updates instead adapt selected existing parameters rather than adding modules. Their expressivity relative to adapters depends on the selected locations and parameter budget. A task-adaptive policy selects those locations using evidence about the current task or user. The resulting update space can target changes that a fixed subset misses, but selection itself consumes resources and does not guarantee that the complete training path fits the device.

The useful savings depend on what the runtime can omit. Unselected parameters need no gradient or optimizer state, but their operations may still lie on the backward path to selected parameters. A policy must therefore compare the measured update cost with its adaptation benefit; the fraction of trainable parameters alone does not predict the reduction in compute. Because layers do not contribute equally to local performance gains, identifying the minimal subset with the highest adaptation value lets the system spend scarce training capacity where it changes the model's behavior most.

11.4.4.1 Sparse update design

Let a neural network be defined by parameters $\theta = \{\theta_1, \theta_2, \ldots, \theta_{N_L}\}$ across N_L layers. In standard fine-tuning, gradients are computed and updates are performed on all parameters:

$$\theta_i \leftarrow \theta_i - \eta \frac{\partial \mathcal{L}}{\partial \theta_i}, \quad \text{for } i = 1, \ldots, N_L$$

In task-adaptive sparse updates, the system selects a small subset $\mathcal{S} \subset \{1, \ldots, N_L\}$ such that only parameters in $\mathcal{S}$ are updated:

$$\theta_i \leftarrow \begin{cases} \theta_i - \eta \frac{\partial \mathcal{L}}{\partial \theta_i}, & \text{if } i \in \mathcal{S} \\ \theta_i, & \text{otherwise} \end{cases}$$

The challenge lies in selecting the optimal subset $\mathcal{S}$ given memory and compute constraints. A principled strategy is contribution analysis, an empirical method that estimates how much each layer contributes to downstream performance improvement. For example, one can measure the marginal gain from updating each layer independently:

1. Freeze the entire model.
2. Unfreeze one candidate layer.
3. Finetune briefly and evaluate improvement in validation accuracy.
4. Rank layers by performance gain per unit cost (e.g., per KB of trainable memory).

This layer-wise profiling yields a candidate ranking from which $\mathcal{S}$ can be constructed subject to a memory budget. Gains measured one layer at a time need not add when layers train together, and shared backward dependencies change the combined cost. The selected set therefore needs a joint quality and resource check, with the profiling cost included in the adaptation budget.

In a complete selective fine-tuning pipeline, that ranking is the middle of three stages (figure 11.9). Stage 1 meta-trains the network in the cloud, which means training its initial weights so that a few later updates on new data adapt it well (the idea behind model-agnostic meta-learning, taken up in section 11.7.3). Stage 2 ranks layers or channel groups by adaptation benefit per unit of device cost and lets the device budget set the cut, and Stage 3 updates only the units above the cut on local data. The backward pass still runs through any frozen layer that lies between the loss and a selected unit, so the selected set must be costed under the complete execution schedule.

A concrete example is TinyTrain, a method designed to allow rapid adaptation on-device (Kwon et al. 2024). TinyTrain combines offline preparation with a task-adaptive sparse-update method that dynamically selects layers or channels to update using user data, memory limits, and compute limits. At runtime, the system updates only the selected subset rather than treating the whole network as trainable.

In implementation, selective layer updating turns the ranking into `requires_grad` decisions or equivalent graph masks. In the profiling-driven version, contribution scores and the memory limit decide which layers are unfrozen (listing 11.3).

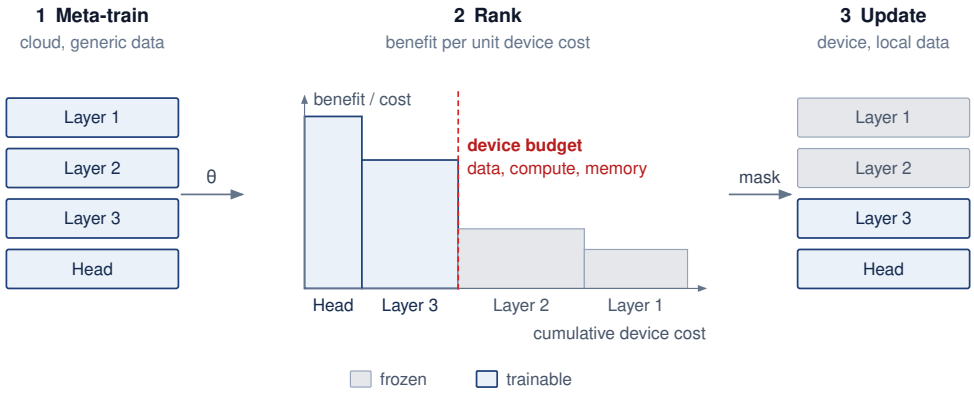

Figure 11.9: Selective On-Device Fine-Tuning: (1) The full network is meta-trained in the cloud on generic data. (2) Each layer or channel group is ranked by its expected adaptation benefit per unit of device cost, and the device's data, compute, and memory budget sets the cut. (3) On the device, only units above the cut are updated on local data while the rest stay frozen. Scores and cut position are schematic.

Listing 11.3: Selective Layer Updating: This technique allows fine-tuning specific layers of a pretrained model while keeping others frozen, optimizing computational resources for targeted improvements.

```
# Selective layer update based on contribution analysis
# Layers selected via profiling: conv2 and fc have highest
# accuracy-per-KB impact for this task

for name, param in model.named_parameters():
    if "conv2" in name or "fc" in name:
        param.requires_grad = True  # Train high-impact layers
    else:
        param.requires_grad = False  # Freeze low-impact layers

# Parameter fraction follows layer shapes; adaptation quality requires evaluation
# Memory savings: gradient storage only for selected layers
```

Consider a user wearing an augmented reality headset that performs real-time object recognition. As lighting and environments shift, the system must adapt to maintain accuracy, but training must occur during brief idle periods or while charging. TinyTrain addresses this constraint with task-adaptive sparse updates, in which the deployment process selects a small update set that fits the device's memory and compute budget. The approach keeps adaptation fast, energy-efficient, and memory-aware when the target task and device profile match the method's assumptions.

Task-adaptive sparse updates still introduce system-level trade-offs. Contribution analysis has a nontrivial computational cost even when it occurs during pretraining or initial profiling, so the deployment pipeline must budget for it. Stability also requires attention: if too few parameters are selected for updating, the model may underfit the target distribution, so teams need validation thresholds for the selected subset before deployment. The selection policy must also account for hardware-specific execution costs because some parameters show high contribution scores but prove expensive to update on certain architectures. Despite these trade-offs, task-adaptive sparse updates provide a practical mechanism to scale adaptation across deployment contexts from microcontrollers to mobile devices (Kwon et al. 2024).

11.4.4.2 Adaptation strategy comparison

The design target is the smallest update space that captures the local change. Full fine-tuning pays every term of equation 11.1. Bias-only updates and head-only training fit minor calibration or new output categories. LoRA and residual adapters add bounded capacity beside a frozen backbone. Task-adaptive sparse updates spend profiling effort to select valuable existing parameters. A larger

update space is justified only if the measured quality gain fits the device's training budget. Each strategy shrinks a different budget term, and the choice must still be checked against expressivity, runtime support, and selection cost. Table 11.4 sets these adaptation strategy trade-offs side by side.

Table 11.4: Adaptation Strategy Trade-Offs: Six approaches to model adaptation compared by trainable parameters, the term of the peak-memory budget (equation 11.1) they shrink, expressivity, suitability, and system requirements. Methods that shrink only the trainable state M_{train} still keep the activations their backward pass needs, while head-only training also removes most of the activation term by stopping the backward pass at the classifier.

Technique	Trainable parameters	Memory overhead	Expressivity	Use case suitability	System requirements
Full Fine-Tuning	All parameters	Highest; every term of the peak budget	Highest	Large domain shift on capable devices	Full backward support and memory headroom
Bias-Only Updates	Bias terms only	Minimal M_{train}; activations remain	Low	Simple personalization; low variance	Extreme memory/compute limits
Head-Only Training	Final classifier	Minimal; backward pass stops at the head	Low	New output classes; custom wake words	Frozen encoder; reusable cached features
LoRA	Low-rank factors per adapted matrix	Low M_{train}; activations remain	Moderate to High	One mergeable adaptation per user	Merge step and requantization check for serving
Residual Adapters	Adapter modules	Low M_{train}; activations remain	Moderate to High	Switching among several contexts on mobile	Mobile-class SoCs with runtime support
Sparse Layer Updates	Selective parameter subsets	Variable; depends on the selected layers	High (task-adaptive)	Real-time adaptation; domain shift	Requires profiling or meta-training

Once an adaptation mechanism fits the device's memory and execution budget, it must still learn from the available examples. Reducing trainable state does not create labels or ensure that a user's recent observations represent future use. The next constraint is therefore statistical: extracting reliable learning signal from sparse local data while retaining useful prior behavior.

11.5 Data Efficiency

A user corrects smartphone keyboard autocorrect perhaps three times a day. The system cannot require 10,000 labeled sentences to train a local language model. Data efficiency at the edge means learning aggressively from a sparse, noisy, and uncurated trickle of implicit user feedback, maximizing the signal extracted from each interaction.

The systems trade-off is data collection cost against adaptation quality. Every retained example costs user friction, energy, storage, and privacy exposure. A voice assistant adapting to speech samples weighs model improvement against battery drain and user comfort with always-on recording, and a keyboard chooses between logging shallow keystroke statistics and keeping full typing sessions. Unlike a central pipeline, the device cannot reach into a larger dataset when its own evidence runs short, and an application that needs a decision before enough evidence accumulates may be better served by a validated fixed model than by immediate learning. Whatever the device does learn must also stay compatible with the adaptation mechanism of section 11.4 and with any federated protocol that later consumes it.

The strategy choice depends on which scarcity binds first, and practical systems usually combine several: few-shot learning handles cases where only a handful of labeled or weakly labeled examples must drive personalization, whereas streaming updates handle observations that arrive incrementally and cannot be batched in advance. Because continual streaming updates can overwrite earlier behavior, experience replay spends memory to retain and reuse scarce examples, while data compression shrinks those stored histories so that replay and lightweight training remain feasible within edge memory budgets. Data scarcity thus becomes an allocation problem, in which the device spends labels, memory, compute, and privacy budget where they produce the most reliable local signal.

11.5.1 Few-shot learning and data streaming

A user enrolling a custom wake word will record five audio samples during setup, not five thousand, and subsequent corrections arrive one interaction at a time during ordinary use. Because local learning must proceed from either a tiny enrollment set or a passive trickle of user interactions and ambient sensor readings, it relies on two complementary adaptation strategies: few-shot learning, in which models generalize from a small, static set of examples, and streaming adaptation, where incremental observations support updates scheduled within device budgets.

Few-shot adaptation uses limited labeled or weakly labeled evidence for a new task or user condition (Wang et al. 2020). Bias-only updates, adapter modules, and prototype-based classification (which labels an input by the nearest per-class average embedding of the few labeled examples) restrict what must be learned, but each still needs validation against overfitting. Let $\mathcal{D}_{\text{shot}} = \{(x_i, y_i)\}_{i=1}^{n_{\text{shot}}}$ denote the local support set; here n_{shot} counts all examples in the set, rather than examples per class. A candidate policy might permit few gradient steps ($n_{\text{steps}} \ll 100$) and a small trainable subset ($|\theta_{\text{updated}}| \ll |\theta|$). These are illustrative restrictions, not universal feasibility conditions. The actual step and parameter budgets follow from measured device cost and the quality target, including retained performance on earlier tasks.

Keyword spotting (KWS) systems are a common target for data-efficient on-device adaptation; Speech Commands provides a dataset for training and evaluating limited-vocabulary KWS models (Warden 2018). A detector must distinguish its target phrase from background speech while meeting latency and false-activation requirements. One design uses a pretrained acoustic encoder, such as a small convolutional or recurrent network, to map audio into embeddings for a lightweight classifier. A shared encoder supplies reusable features, while local examples can support a custom wake word or an underrepresented accent. Whether those examples justify local training depends on their coverage, transfer restrictions, and the improvement over the shared detector.

Few-shot adaptation can train an output classifier or another small parameter subset from utterances collected directly on the device. A user might provide 5–10 recordings of a custom wake word while the acoustic encoder remains frozen. Those positive examples define what to recognize, but the false-activation rate on other speech and background noise is a separate measurement, and the personalized detector should not be accepted until it has been made.

Head-only adaptation is cheap in a way the peak-memory budget of equation 11.1 makes precise. Because the backward pass stops at the classifier, the encoder's saved activations leave the activation term, and the encoder outputs for the support set can be computed once and reused across update steps. What remains is running the encoder forward, holding the support examples or their features, and evaluating the adapted classifier. Those costs, not the size of the head, decide whether the update fits a particular phone or microcontroller.

Many on-device scenarios instead call for streaming adaptation, in which the model learns incrementally as data arrives over time rather than from a fixed support set (Hayes et al. 2020). Let $\{x_t\}_{t=1}^{\infty}$ represent the observations selected from the stream for admitted updates. In streaming settings, observations arrive continuously, but the runtime admits training only when memory, energy, and foreground-latency budgets permit. The following rule describes one admitted update; it does not require a gradient step on every arrival. The index t counts these update steps. The computable loss may come from implicit feedback, pseudo-labels, self-supervised objectives, reconstruction error, or a task-specific proxy rather than explicit human labels. The model update can be written generically as:

$$\theta_{t+1} = \theta_t - \eta_t \nabla \mathcal{L}(x_t; \theta_t)$$

where η_t is the learning rate at time t. This form of adaptation is sensitive to noise and drift in the input distribution, and thus often incorporates mechanisms such as learning rate decay, meta-learned initialization, or update gating to improve stability.

The same strategies apply beyond keyword spotting. In wearable health devices, a model that classifies physical activities may begin with a generic classifier and adapt to user-specific motion patterns using only a few labeled activity segments. In smart assistants, user voice profiles are fine-tuned over time using ongoing speech input, even when explicit supervision is unavailable. Local corrections, repetitions, or downstream task completion can provide implicit signals, but only after the system defines which event counts as feedback and tests its validity. A user may accept a suggestion because it is convenient rather than correct, abandon a task for reasons unrelated to the

model, or correct an error that the model itself induced. Production adaptation should therefore separate observed events from inferred labels, retain a nonadapting control cohort, measure delayed outcomes when available, and gate learning when the proxy-to-outcome relationship drifts. A stream of such updates also raises a stability problem that few-shot and streaming rules do not address on their own.

11.5.2 Experience replay

When a streaming update rule trains only on a user's newest interactions, gradient steps on recent examples overwrite the weights that recognized earlier classes, causing accuracy on prior tasks to collapse within a few updates. Retaining a small sample of earlier observations on the device reduces this failure mode without eliminating it.

Definition 11.3: Catastrophic forgetting

Catastrophic forgetting is the phenomenon in which a neural network trained sequentially on new tasks or data distributions rapidly loses performance on previously learned tasks, because gradient updates that optimize the new objective overwrite weight configurations that supported the old ones.

1. **Significance**: Without explicit mitigation, fine-tuning a deployed classifier on a new domain can sharply reduce prior-task accuracy, even when the new domain is closely related to the old one (Kirkpatrick et al. 2017). This is the central failure mode of on-device and federated learning. Every device that adapts locally to its user's recent data risks erasing the general capabilities the centralized training run paid for, and the loss is invisible until inference encounters out-of-recent-context inputs.
2. **Distinction**: Unlike overfitting (a generalization failure on a single fixed task) and distribution shift (a deployment-time mismatch between training and inference data), forgetting is a training-time phenomenon caused by the sequential structure of the update stream itself. Updates for new data can interfere with earlier solutions even when a representation capable of supporting both tasks exists.
3. **Common pitfall**: A frequent misconception is that a small experience-replay buffer is sufficient protection. A small buffer can hold aggregate accuracy on a held-out set while rare classes, which it samples too seldom, continue to degrade. Forgetting must therefore be measured per class or per cohort, and the remedies range from capacity-proportional or prioritized buffers to weight-regularization methods such as elastic weight consolidation, which store no examples (section 11.8.2).

Experience replay addresses forgetting by retaining examples from earlier learning episodes and reusing them during later updates. Replay also appears in reinforcement learning (Mnih et al. 2015); here its role is to keep earlier local observations represented as new data arrives. It is one stability mechanism with an explicit storage cost. Section 11.8.2 returns to the problem with parameter-importance penalties, bounded update footprints, and rollback.

Unlike server-side replay strategies that rely on large datasets and extensive compute, on-device replay must operate with extremely limited capacity, often with tens or hundreds of samples, and must avoid interfering with user experience (Rolnick et al. 2019). Buffers may store only compressed features or distilled summaries, and updates must occur opportunistically (e.g., during idle cycles or charging). These system-level constraints reshape how replay is implemented and evaluated in the context of embedded ML.

Let $\mathcal{M}$ represent a memory buffer that retains a fixed-size subset of training examples. At time step t, the model receives a new data point (x_t, y_t) and appends it to $\mathcal{M}$. A replay-based update then samples a batch $\{(x_i, y_i)\}_{i=1}^{B}$ of size B from $\mathcal{M}$ and applies a gradient step:

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta \left[\frac{1}{B} \sum_{i=1}^{B} \mathcal{L}(x_i, y_i; \theta_t) \right]$$

This is the streaming update of section 11.5.1 with one change of data source, from the newest observation to a batch drawn from $\mathcal{M}$. The systems cost lies in what that change requires, namely buffer storage that competes with model state and forward and backward work on old examples in every admitted window. Over time, this replay mechanism allows the model to reinforce prior knowledge while incorporating new information.

A practical on-device implementation might use a ring buffer to store a small set of compressed feature vectors rather than full input examples. The buffer has a fixed capacity and overwrites its oldest entry once full, so one parameter bounds its memory and sets the balance between retained history and new data (listing 11.4).

Listing 11.4: Replay Buffer: Implements a circular storage mechanism for efficient memory management in constrained environments. This approach allows models to efficiently retain and sample from recent data points, balancing the need to use historical information while incorporating new insights.

```
import random


# Ring buffer for experience replay on memory-constrained devices
# Stores compressed features (not raw inputs) to minimize footprint
class ReplayBuffer:
    def __init__(self, capacity):
        self.capacity = capacity  # Fixed size, e.g., 1000 examples
        self.buffer = []
        self.index = 0  # Circular write pointer

    def store(self, feature_vec, label):
        if len(self.buffer) < self.capacity:
            self.buffer.append((feature_vec, label))  # Fill phase
        else:
            self.buffer[self.index] = (
                feature_vec,
                label,
            )  # Overwrite oldest
        self.index = (self.index + 1) % self.capacity  # Wrap around

    def sample(self, k):
        # Uniform sampling within the retained buffer; older examples are absent
        return random.sample(self.buffer, min(k, len(self.buffer)))
```

This implementation maintains a fixed-capacity cyclic buffer, storing compressed representations (e.g., last-layer embeddings) and associated labels. Uniform sampling here covers only the retained recent window; it does not remove the history bias introduced by overwriting old examples. Reservoir sampling, which keeps each new arrival with a probability that shrinks as the stream grows so that every observation seen so far is equally likely to remain in the buffer, is an alternative when the goal is to represent the full stream, though any finite buffer still requires coverage checks.

In TinyML applications, experience replay has been applied to problems such as gesture recognition, where devices must continuously improve predictions while observing a small number of events per day. Instead of training directly on the streaming data, the device stores representative feature vectors from recent gestures and uses them to finetune classification boundaries periodically. Similarly, in on-device keyword spotting, replaying past utterances can improve wake-word detection accuracy without the need to transmit audio data off-device.

Experience replay spends storage and additional update work to preserve earlier behavior. Raw inputs can exceed storage or retention budgets, especially in vision and audio applications. Cached feature vectors reduce that footprint, but an update starting from those vectors trains the downstream model without reconstructing the original encoder computation. If the encoder changes, old features may no longer match its current representation; the system must retain a compatible encoder, refresh the features from retained inputs, or expire incompatible entries. Persistent replay also consumes flash write cycles. Buffer contents, representation versions, and replay frequency therefore belong in the same resource and retention policy.

11.5.3 Data compression

Experience replay is limited by how many examples fit in the buffer: storing raw 224×224 camera frames or uncompressed audio waveforms inside an on-device replay buffer fills a wearable's or microcontroller's memory after only a few dozen observations. When raw training inputs are too large, noisy, or redundant to store and process directly, transforming them into lower-dimensional embeddings or compact encodings preserves task-relevant structure while shrinking memory and compute costs. Compressed representations serve two complementary goals. They reduce the footprint of stored data, allowing devices to maintain longer histories or replay buffers under tight memory budgets (Hayes et al. 2020). They also simplify the learning task by projecting raw inputs into more structured feature spaces, often learned via pretraining or meta-learning, in which efficient adaptation is possible with minimal supervision.

One common approach is to encode data points using a pretrained feature extractor and discard the original high-dimensional input. For example, an image x_i might be passed through a CNN to produce an embedding vector $\mathbf{z}_i = f(x_i)$, where $f(\cdot)$ is a fixed feature encoder. This embedding captures visual structure (e.g., shape, texture, or spatial layout) in a compact representation, usually ranging from 64 to 512 dimensions, suitable for lightweight downstream adaptation.

Mathematically, training can proceed over compressed samples $(\mathbf{z}_i, y_i)$ using a lightweight classification or projection head. Let θ represent the trainable parameters of this head, which is typically a small neural network that maps from compressed representations to output predictions. As each example is presented, the model parameters are updated using gradient descent:

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta \mathcal{L}(g(\mathbf{z}_i;\theta), y_i)$$

The rule is the same gradient step once more, now applied to stored embeddings $\mathbf{z}_i$ and a small head whose prediction is $g(\mathbf{z}_i;\theta)$, rather than to raw inputs. Only the head's parameters θ change, so its memory is the head-only case of the budget in equation 11.1, and the encoder runs once per example when the embedding is stored rather than on every update. This bounds the trainable portion of the model, but feasibility still depends on the head's activations and optimizer state alongside the replay buffer and runtime reserves.

When replay storage rather than model compute binds, a device can go further and learn a dictionary of recurring sensor-trace patterns, storing each example as a few nonzero coefficients over that dictionary instead of as raw samples. A dataset of traces $\mathbf{X}$ then factorizes as $\mathbf{X} \approx \Phi_{\text{dict}}\mathbf{C}$, where the columns of Φ_{dict} are the basis patterns and the block-sparse coefficient matrix $\mathbf{C}$ records which patterns are active in each example. The saving is real only if the dictionary plus all stored coefficients is smaller than the raw replay examples. Updating the dictionary changes the representation of every stored example, so reconstruction quality and compatibility become part of the replay policy.

Compressed representations can reduce how long raw inputs remain stored, but an embedding is not anonymized merely because it is lower dimensional. It may preserve identity, attributes, or enough structure for inversion, and it must receive the same access-control, retention, and deletion analysis as the source data unless a separate privacy mechanism establishes a stronger guarantee. Compression can also act as an implicit regularizer, smoothing learning and mitigating overfitting when only a few training examples are available.

In keyword spotting, raw audio can first be transformed into mel-frequency cepstral coefficients (MFCCs),[12] a compact, lossy representation derived from the speech power spectrum. Storing these vectors instead of waveforms reduces the replay payload. The complete adaptation budget must still include feature extraction, model state, and the working memory needed for updates.

[12] **MFCCs (mel-frequency cepstral coefficients)**: For a 320-sample audio window (20 ms at 16 kHz), 16-bit samples occupy 640 bytes, while 12–13 coefficients stored at 32 bits each occupy approximately 50 bytes. The roughly $13\times$ reduction compares one waveform window with one feature vector; overlapping windows, additional features, and metadata change the storage ratio for a complete recording. Replay capacity depends on that complete representation, and keyword detection quality must be evaluated separately from byte savings.

11.5.4 Data efficiency strategy comparison

Few-shot adaptation, replay, and compression each spend a different scarce resource and fail in a different way. Few-shot adaptation spends scarce labels on a small informative support set, and weak pretrained features can defeat it. Replay spends memory and compute to preserve prior behavior; raw observations can overflow device storage when memory is the bottleneck, and a buffer can omit rare cases. Compressed representations extend the history that fits on the device, but compression can discard task-relevant information, and fixed embeddings freeze the wrong representation when the feature extractor itself needs to adapt. Combining mechanisms does not escape these on-device

learning trade-offs (table 11.5), because every combination must fit the same memory, energy, and privacy budget.

Table 11.5: On-Device Learning Trade-Offs: Technique choice depends on the binding edge constraint: few-shot adaptation minimizes labeled-data needs, replay improves stability under drift at the cost of memory, and compressed representations extend retention when raw data is too large or too sensitive to keep.

Technique	Data requirements	Memory/compute overhead	Use case fit
Few-Shot Adaptation	Small labeled set (n_{shot} examples)	Low	Personalization, quick on-device finetuning
Experience Replay	Streaming data	Moderate (buffer & update)	Non-stationary data, stability under drift
Compressed Representations	Unlabeled or encoded data	Low to Moderate	Memory-limited devices; privacy-sensitive contexts

In practice, these methods are not mutually exclusive. Many real-world systems combine them to achieve robust, efficient adaptation. For example, a keyword spotting system may use compressed audio features (e.g., MFCCs), finetune a few parameters from a small support set, and maintain a replay buffer of past embeddings for continual refinement.

Efficient adaptation and data strategies establish how a single device can learn from its user. Sharing those gains requires another decision: whether local updates contain information useful to a common model and can be combined within communication and privacy constraints. Federated learning supplies that coordination protocol. Personal state can remain local when sharing would not improve the population objective, and privacy controls must separately limit what transmitted updates disclose.

11.6 Federated Learning Algorithms

The voice assistant of the chapter's opening, now deployed to millions of homes whose intermittently available devices contribute updates when local conditions permit, shows what local adaptation alone cannot do. One household pronounces "data" as "DAY-tuh" and another as "DA-tuh"; a third, a technical household, says "machine learning" every day, while a fourth never has. Local adaptation can specialize each model to its own household, but nothing it learns travels. If the fourth household starts using the phrase, its model must learn it from scratch, even though millions of other devices already handle it. Federation offers that transfer when the observations support a common objective. The system must decide which changes improve shared representations and which should remain personal, then evaluate whether aggregation helps the intended users.

Federated learning (definition 11.2) supplies that transfer while each device keeps its raw examples local. The approach turns data locality into one privacy control, not a complete privacy guarantee. A system must still treat gradients, model deltas, participation timing, device identifiers, and aggregate statistics as sensitive outputs because these signals can reveal information about local users (Kairouz and McMahan 2021; Zhu et al. 2019). Of the three architectures in figure 11.2, only federated learning pools experience across users without moving raw examples, and it buys that position with a protocol.

11.6.1 Cross-device and cross-silo federations

Before choosing a federated optimization or privacy protocol, an architect must specify who the participants are, because a mobile cohort-sampling protocol applied to ten collaborating hospitals breaks scheduling, weighting, and fault recovery, and a consortium protocol that expects stable round membership fails on millions of intermittently connected phones. The voice assistant is cross-device, with millions of unreliable phones, whereas a group of hospitals that want to train a complication-prediction model without centralizing patient records is cross-silo, with a few reliable institutions that each hold a large dataset. That distinction changes the failure model, scheduling policy, and privacy unit before any optimizer is selected (table 11.6) (Kairouz and McMahan 2021).

Table 11.6: Cross-Device and Cross-Silo Federated Learning: The participant population determines scheduling, trust, weighting, and privacy semantics. A protocol designed for millions of unreliable phones is not automatically appropriate for a consortium of hospitals.

Design dimension	Cross-device federation	Cross-silo federation
Participants	Large population of intermittently available personal devices	Small set of known organizations or sites
Round semantics	Sample an eligible cohort; expect stragglers and dropouts	Coordinate a stable membership; investigate missing sites
Trust boundary	Treat client software and local data as potentially compromised	Authenticate institutions, but still isolate each data custodian
Weighting default	Decide whether the objective represents examples or users	Weighting follows the agreed institutional objective, not site size by default
Privacy unit	Usually a user or device across its participation history	Usually a person or record within a site's data, sometimes the site itself

11.6.2 Learning protocols

Once the participant population is defined, a federated round must coordinate devices that hold different data distributions, connect over constrained links, and drop out unpredictably. If devices trained on unaligned starting versions and uploaded raw weight tensors whenever they came online, the coordinator would sum incompatible parameter spaces and stall on offline stragglers. Practical distributed coordination protocols must therefore solve three engineering challenges simultaneously. Local training must produce compatible updates despite non-IID data distributions. Aggregation must combine those updates efficiently despite bandwidth constraints. Timing must be coordinated across devices with heterogeneous availability patterns.

11.6.2.1 Local training

Local training turns the device's adaptation policy into a contribution to a shared model. Every contribution needs a known starting version, a compatible parameter schema, and a declared local update rule. Otherwise, arrays with the same shape may describe changes to different models. A synchronous protocol therefore starts every participating client from the same shared version, trains on local examples under that rule, and returns the difference from the starting parameters, while personal parameters stay outside the aggregate.

The process repeats with a new shared version for each accepted round. Update frequency depends on device availability, local work, and communication cost. A secure transmission channel protects the transfer, but the receiving server can still inspect the update unless the aggregation protocol separately hides individual contributions. Model compatibility and update confidentiality are therefore distinct protocol obligations.

11.6.2.2 Federated aggregation protocols

The most widely used rule for combining client updates is federated averaging (FedAvg), which has become a canonical algorithm for federated learning (McMahan et al. 2017). In FedAvg, each device trains its local copy of the model using SGD on its private data. The algorithm is local SGD, which synchronizes only after several local steps to save communication, run with two changes that the data center version never faces. Only a sampled subset of clients participates in each round, and each client's data is a non-IID shard of the population. The local steps that save communication therefore also pull each client toward its own optimum.

↩ Local SGD and the trade-off between synchronization frequency and convergence that it exposes are established in section 6.4.

Formally, let $\mathcal{D}_k$ denote the local dataset on client k, and let θ_k^t be the parameters of the model on client k at round t. Each client performs E epochs of SGD on its local data, yielding an update θ_k^{t+1}. The central server then aggregates these updates as:

$$\theta^{t+1} = \sum_{k=1}^{N_{\text{clients}}} \frac{n_k}{n_{\text{round}}} \theta_k^{t+1} \tag{11.3}$$

where $n_k = |\mathcal{D}_k|$ is the number of samples on device k, $n_{\text{round}} = \sum_k n_k$ is the total number of samples across participating clients, and N_{clients} is the number of active devices in the current round. Because

every client starts from the same shared model and the weights sum to one, averaging the returned models is equivalent to adding the weighted average of the returned differences to the starting model, which is the form the local-training protocol transmits. Bare E counts local epochs; the energy quantities of equation 11.2 always carry a subscript.

This formula optimizes an *example-weighted* objective, in which a client with ten times as many local examples receives ten times the influence. That choice is appropriate when every example should contribute equally, but it can let prolific users or large institutions dominate the model. A *user-weighted* objective instead assigns each participating client weight $1/N_{\text{clients}}$, optionally capping or normalizing each client's local contribution before aggregation. Neither choice is universally fair. The system must declare whether its unit of interest is an example, a user, an organization, or a deployment cohort, then evaluate that population rather than inherit n_k/n_{round} as an implementation default.

A deployed round wraps the averaging rule in policy that the equation leaves unstated (figure 11.10). Equation 11.3 fixes how completed models combine, but not which clients are invited, how long the coordinator waits, or how few completions force a round to be abandoned. Those protocol choices decide which contributions enter the sum. A client that misses the deadline contributes nothing, and its examples are absent from n_{round} for that round. Over-selection and its population-coverage trade-off are developed in section 11.7.1.

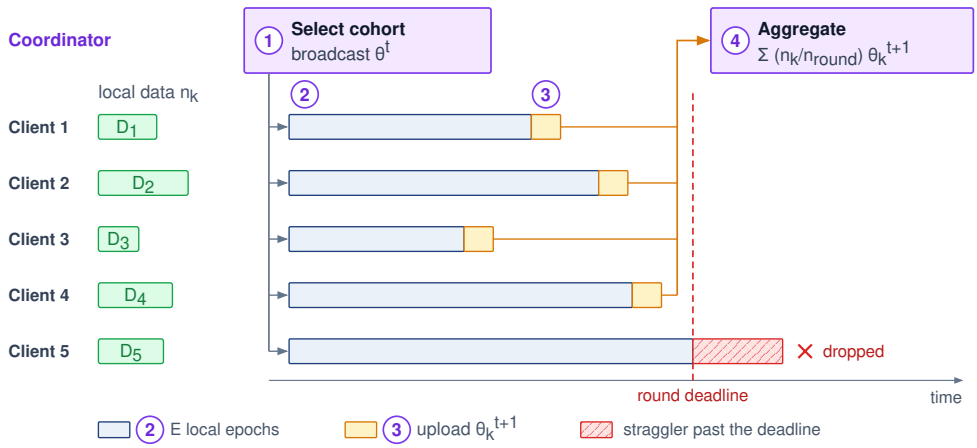

Figure 11.10: Federated Averaging Protocol: One FedAvg round in four steps. (1) The coordinator selects a cohort, over-selecting to absorb dropouts, and broadcasts the global model θ^t. (2) Each client runs E local epochs on its private data. (3) Clients that finish before the round deadline upload θ_k^{t+1}; a straggler that misses the deadline is dropped. (4) The coordinator forms the example-weighted average of equation 11.3, which becomes the next global model.

This basic structure introduces a number of design choices and trade-offs. The number of local epochs E sets the balance between computation and communication. A larger E reduces communication frequency but risks divergence if local data distributions vary too much. The selection of participating clients affects convergence stability and fairness. In real-world deployments, not all devices are available at all times, and hardware capabilities may differ substantially, requiring robust participation scheduling and failure tolerance.

11.6.2.3 Time to target under client drift

A federated deployment must reach a quality target within its time, energy, and participation budgets. More local training can reduce how often clients communicate, but it also changes their updates and the time each selected cohort takes to finish. The engineering comparison is therefore time to an agreed quality target, not communication rounds alone.

⇄ **Contrast:** Federated rounds over intermittent wide-area clients contrast with the synchronous AllReduce collectives on low-latency cluster fabrics established in section 5.6.

What the FedAvg controls count FedAvg exposes the number of participating clients N_{clients}, local epochs E, and communication rounds R as controls (McMahan et al. 2017). Their product $N_{\text{clients}}ER$

counts client-epochs when participation is fixed, which makes it a tally of the work a run performs rather than a measure of its progress. Clients can have different dataset sizes, and equal client-epoch counts can process different examples or produce conflicting updates, so two runs with the same $N_{\text{clients}}ER$ can end at very different quality. A convergence result is useful only when its assumptions match the optimizer, client sampling, and data distributions.

Non-IID data impact Non-IID data is where federated averaging departs from the data center version of local SGD. Each local step moves a client toward the optimum of its own shard. After many local steps the client models sit near different optima, and their weighted average lands at a point that need not minimize the population objective. This client drift grows with the number of local steps and with how far the client distributions differ, and it degrades federated averaging relative to i.i.d. baselines (T. Li et al. 2020; Zhao et al. 2018). It therefore eats into the benefit of fewer communications exactly where communication is most expensive. How strongly depends on the partition, update rule, participation policy, and local step size. No universal heterogeneity score predicts a round count across applications, which is why section 11.6.2.3 measures the trade-off rather than computing it.

When statistical heterogeneity causes FedAvg to diverge or stall, two foundational algorithmic extensions address the root cause of client drift:

1. **FedProx** (T. Li et al. 2020) adds a proximal penalty proportional to $\|\theta - \theta^t_{\text{global}}\|^2$ to each client's local loss function. This penalty constrains local updates to remain in the neighborhood of the current global model θ^t_{global}, stabilizing convergence under statistical heterogeneity and permitting devices to execute variable numbers of local epochs without destabilizing the server aggregate.
2. **SCAFFOLD** (Kairouz and McMahan 2021) corrects the local update direction directly using control variates. By maintaining a server control variate and per-client control variates that estimate the difference between client-local and global gradient directions, SCAFFOLD applies an additive correction to each local SGD step. Each step subtracts the client's own control variate and adds the server's, so the client moves along the estimated global direction rather than the direction its local data favors. This correction counters client drift under heterogeneous local data, although it roughly doubles the communication volume per round by transmitting control variates alongside model parameters. Each client's control variate is useful only when that client returns, so the method suits cross-silo federations better than cross-device rounds that rarely sample the same device twice.

Measuring the local-work trade-off Compare candidate local-epoch settings under the same initialization, data partition, participation policy, and held-out quality target. For each setting E, record the first validated round $R(E)$ that reaches the target and the elapsed duration $t_r(E)$ of each completed round, including client selection, local work, transfer, aggregation, and any waiting. The observed time to target is

$$T_{\text{target}}(E) = \sum_{r=1}^{R(E)} t_r(E).$$

Each $t_r(E)$ is one instance of the round budget sketched at the opening of the chapter. Read through the Fleet Law, it is the local training on the slowest client the round still awaits, plus the upload, plus the wait for enough reports, less whatever overlaps. The communication piece is easy to estimate. Over a constrained mobile uplink of 1–10 Mb/s, the chapter's reference adapter update of 50 MB takes about 7 minutes to upload at the low end and about 42 seconds at the high end. That cost recurs in every one of the $R(E)$ rounds. A setting of E that halves $R(E)$ therefore removes half of the cumulative upload time. It pays off only if the extra local training it adds to each round, on the slowest awaited client, costs less than the uploads it removes. Client drift keeps that comparison from being settled on paper, because more local steps can also raise $R(E)$, and only measured runs can rank the settings.

A useful experiment log records local epochs, completed clients, rounds to target, elapsed time, transmitted bytes, and device energy. If a run never reaches the quality target within the budget, report that outcome rather than extrapolating an invented convergence curve. Repeated runs and

cohort-level validation distinguish a stable improvement from a fortunate sample of participating devices.

Communication-computation trade-off Figure 11.11 draws the two opposing paths of local work E as a timeline of rounds. Each round's width is its duration $t_r(E)$, each row's round count is $R(E)$, and each row ends at its time to target $T_{\text{target}}(E)$. Drift adds a round to the large-E row, and the moderate row finishes first. The large-E row needs half the rounds of the small-E row yet finishes last, so a setting that requires fewer rounds can still prolong total wall-clock duration. A setting remains admissible only if its total consumption also fits within client energy and fleet-coverage budgets.

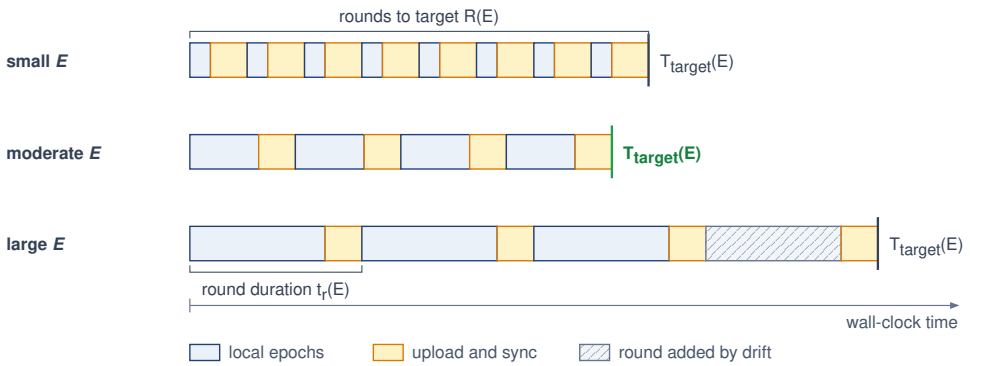

Figure 11.11: Measuring the Local-Work Trade-Off: The local-work setting E acts through two opposing paths. More local epochs lengthen each round, $t_r(E)$, while reducing how often clients synchronize, which can lower the rounds to target, $R(E)$, unless client drift raises it. Only their combination, the measured time to target $T_{\text{target}}(E)$, ranks candidate settings, and a setting is admissible only if it also fits the client energy and population-coverage budgets.

The operating point depends on the deployment. Larger local updates can amortize transfer overhead, but they also change drift, completion time, and which clients finish. The network alone can move the optimum, as a cost model in relative units shows (figure 11.12). Its minimum E^* shifts from $E = 10$ on the fast-network curve to $E = 40$ on the slow-network curve because amortizing expensive transfers justifies more local work. These curves are not measured convergence results. A deployment must establish whether additional local work actually reduces time to target after drift and client completion are counted, and adaptive epoch policies need the same outcome and coverage checks as fixed settings.

Diagnosing a missed target Federated convergence results depend on the update rule, objective, sampling policy, and assumptions about client data (Kairouz and McMahan 2021). Heterogeneity, participation, and local sample counts are therefore inputs to a feasibility assessment, not universal pass/fail thresholds. A deployment must demonstrate that its chosen configuration reaches the declared quality target within time and device budgets, with evaluation that represents the intended population.

When a configuration misses that target, the measured failure guides the next experiment. Conflicting local updates motivate testing less local work, adding proximal regularization (FedProx), incorporating control variates (SCAFFOLD), or switching to personalization. Long waits motivate scheduling or asynchronous alternatives, subject to privacy-protocol compatibility. Sparse local evidence motivates bounded adaptation and stronger validation. None of these changes replaces observations from absent cohorts; each must be reevaluated for quality, resource cost, and population coverage.

The aggregation rule specifies how compatible local models produce a shared model. The measured time to target establishes whether a configuration reaches its objective within budget. A controlled lab can hold participant availability and network conditions steady. A phone fleet cannot, so the implementation must define how selection, deadlines, and failures change the set of updates

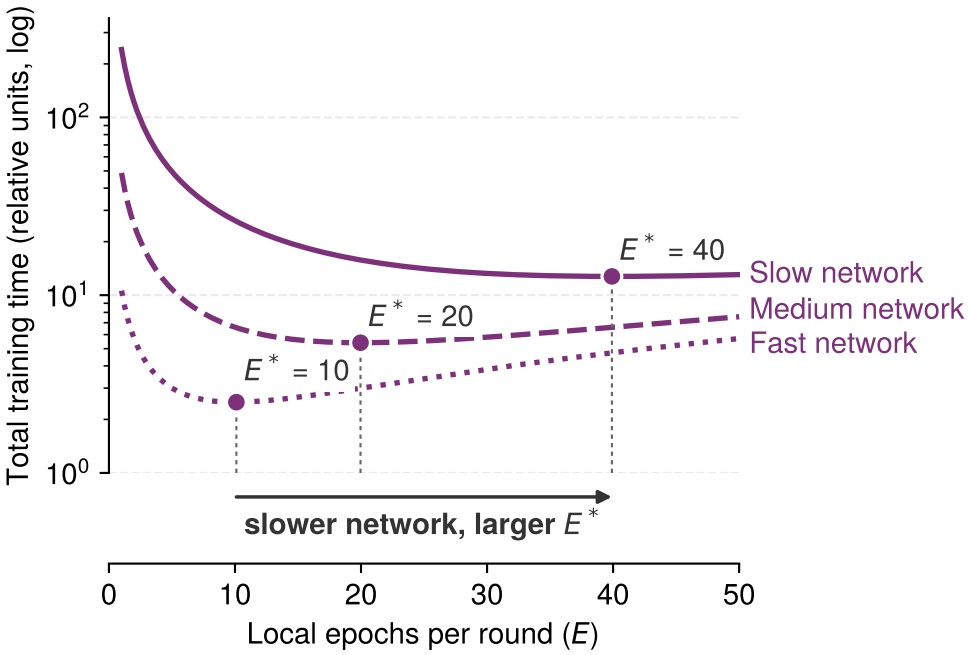

Figure 11.12: Communication-Computation Trade-Off in Federated Learning: As network bandwidth decreases from fast to slow, the time-minimizing number of local epochs shifts rightward because additional local computation amortizes more communication. The curves come from an illustrative cost model in relative units, not measurements; its growing local-work term stands in for the longer rounds and added drift that make excessive local work counterproductive.

that actually enters each aggregate. These protocol decisions connect learning behavior to fleet orchestration.

11.7 Federated Learning Systems at Scale

A round selects a cohort from the deployment population, but some selected devices overheat, lose Wi-Fi, enter power-saving mode, or power off before returning an update. A fleet of millions of smartphones need not participate in one round for these failures to matter. Turning the FedAvg protocol into a production fleet system requires deciding which clients are eligible and how long to wait, compressing updates to fit constrained uplinks, separating personal state from shared weights, enforcing cryptographic and statistical privacy bounds, and reconciling asynchronous updates across heterogeneous hardware.

11.7.1 Client scheduling

The first decision in any round is which devices to invite. Sampling devices uniformly at random without checking their power and network state would interrupt users mid-task, whereas accepting updates only from plugged-in flagship phones on fast Wi-Fi trains a global model skewed toward affluent, overnight-charging cohorts. Because real-world client availability shifts continuously with battery state, charging habits, and network coverage, client scheduling governs both the completion speed and the population coverage of distributed learning.

At a baseline level, federated ML systems define eligibility criteria for participation. A mobile deployment may require devices to be plugged in, connected to Wi-Fi, and idle to limit interference with foreground work and battery use. The chosen criteria determine which subset of the total population is considered "available" for a training round; they are deployment policy rather than requirements of federated learning itself.

Beyond these operational filters, devices also differ in their hardware capabilities, data availability, and network conditions. Some smartphones contain many recent examples relevant to the current

task, while others have outdated or irrelevant data. Network bandwidth and upload speed may vary widely depending on geography and carrier infrastructure. As a result, uniform random selection from the eligible pool does not guarantee coverage of the full deployment population; eligibility has already excluded some devices and their data.

Availability-driven selection introduces participation bias, the fleet-level consequence of the admission test in equation 11.2 (principle 20). Clients with favorable conditions, such as frequent charging, high-end hardware, and consistent connectivity, pass admission and participate repeatedly, while others are systematically underrepresented. This can skew the resulting model toward behaviors and preferences of a privileged subset of the population, raising both fairness and generalization concerns.

Participation bias becomes severe when a small eligible cohort supplies most completed updates while other cohorts rarely participate. The skew is measured per cohort as a coverage ratio, the cohort's share of completed updates divided by its share of the deployed population. A ratio near one means the cohort is represented in proportion to its size, and a ratio well below one means the shared model learns that cohort's behavior at less than its population weight. Completion logs supply the numerator and the deployment inventory the denominator, so the ratio can be computed without inspecting any update. The admission test and the round deadline both remove constrained devices faster than capable ones. Uniform selection among admitted devices passes the resulting skew through unchanged, so a budget cohort can fall from a fifth of the fleet to a sliver of the completed updates (figure 11.13).

Low coverage also feeds itself, because models become more strongly optimized for users with high-end devices and stable connectivity, potentially degrading performance for resource-constrained users who need adaptation the most. A keyboard prediction model might become biased toward the typing patterns of users with flagship phones who charge overnight, missing important linguistic variations from users with budget devices or irregular charging patterns.

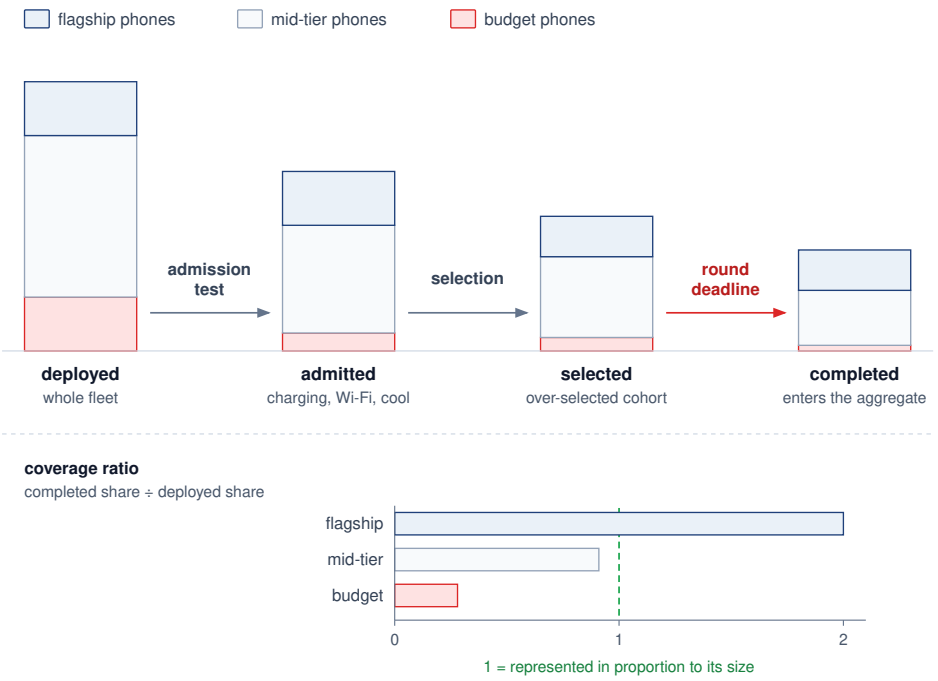

Figure 11.13: From Deployed Fleet to Aggregate: Each stage between the deployed fleet and the aggregate filters devices, and the admission test and round deadline fall hardest on constrained cohorts. The admission test removes most budget phones, uniform selection among admitted devices preserves the skew it inherits, and the round deadline drops the slowest survivors. The coverage ratio, each cohort's share of completed updates divided by its share of the deployed fleet, sizes the result: above 1 for flagship phones, below 1 for budget phones. Cohort shares and stage sizes are illustrative, not measured.

The order of these decisions limits what any remedy can do. Fleet orchestration admits only feasible allocations before a fairness policy divides them (principle 15), and a federated scheduler inherits the same order with a sharper consequence. Its eligibility filter fixes the candidate pool, so a fairness policy can redistribute rounds only among devices the admission test already accepted; a device that never qualifies is invisible to it.

Scheduling must therefore balance completion speed with population coverage. Stratified or quota-based sampling can prioritize underrepresented groups among eligible devices, while cumulative participation records can guide later invitations. Weighting methods require estimates of the population and selection process; they cannot recover observations from a cohort that never participates. Buffered asynchronous aggregation (Nguyen et al. 2021), whose mechanism section 11.7.5.1 describes, staleness-aware methods (Rodio and Neglia 2024), and fairness-oriented weighting (Ma et al. 2024) offer additional controls. Their effectiveness must be checked through completed participation and cohort quality, rather than inferred from the scheduling rule alone. Example weighting in the aggregate (equation 11.3) also gives large local datasets more influence whether or not they represent the target population any better.

Scheduling must also bound how long a round waits on its slowest participants. Network latency, battery depletion, or background process contention can make some devices (stragglers) take $10\times$ longer than the median, and the round cannot aggregate until enough updates arrive.

To reduce waiting for stragglers, production federated learning (FL) systems employ **over-selection**. The server selects a candidate pool size $K_{\text{candidates}}$ larger than the target number of updates K_{target} (the production system of Bonawitz et al. (2019) selects about 130 percent of the target, $K_{\text{candidates}} \approx 1.3 \times K_{\text{target}}$). When enough valid reports arrive, the server can form an aggregate from K_{target} completed updates rather than wait for every selected device. Completion then depends on how quickly the target number of valid reports arrives, but over-selection alone does not guarantee that the target is reached. A round deadline and an explicit acceptance or abandonment rule handle insufficient completion. Secure aggregation adds a participation floor and dropout recovery (section 11.7.4); selecting the fastest reports must remain compatible with that protocol. Accepting the first reports to arrive discards those of the slowest selected devices, which are disproportionately the constrained cohorts that admission already thinned, so over-selection with a round deadline can concentrate the participation bias just described. Shorter rounds must therefore be evaluated against cohort coverage and wall-clock time to the declared quality target. Because these scheduling decisions set convergence rate, model generalization, energy consumption, and user experience together, poor scheduling produces excessive stragglers, overfitting to narrow client segments, or wasted computation.

Enough valid reports let the round close before its slowest clients finish.

11.7.2 Bandwidth-aware update compression

Many of the stragglers that miss a round deadline fail on the network transfer rather than on local computation: transmitting full-precision weights or gradients over an asymmetric mobile uplink can exhaust a client's connectivity window or battery allowance before its update ever reaches the coordinator.[13] Update compression therefore begins with a bottleneck diagnosis: when transfer is limiting, the design decision is which information to quantize, sparsify, or keep local. The smaller payload must then be evaluated against encoding cost, update error, and the rounds required to reach the quality target.

The compression decision separates into three levers, each changing a different part of the communication bottleneck: model compression reduces bits per update, selective sharing reduces which parameters travel, and architectural partitioning keeps some state local. Model compression methods reduce the wire footprint of transmitted updates through quantization, sparsification, or subsampling. Instead of sending full-precision gradients, a client transmits 8-bit quantized updates or communicates only the top-K_{top} gradient elements[14] with highest magnitude, the same encodings that cluster training applies to its collective payloads. At the edge they bind in a different place. They act on a slow, asymmetric uplink rather than a fabric, under secure aggregation that must still combine the encoded updates, and inside a per-round communication total (equation 11.4) that also counts downloads and retries. These transformations are communication mechanisms, not privacy mechanisms; a smaller embedding, sparse update, or low-precision gradient may still expose membership, attributes, or reconstructable features.

13 **Wireless upload asymmetry**: Federated updates use the uplink, whose available throughput can differ substantially from the downlink and vary with coverage and contention. Transfer time depends on encoded update size and effective uplink goodput; energy also depends on radio state and transfer duration. Low-power links may add duty-cycle constraints. Compression is useful when these costs bind the update budget, but its necessity depends on the payload, link, and permitted synchronization interval.

14 **Gradient sparsification**: Selecting the largest 1–10 percent of gradient elements by magnitude reduces transmitted value count by 10–100$\times$ before coordinate indices and other metadata are included. Sparsified SGD and deep gradient compression retain untransmitted residuals for later updates rather than assuming that small components can be discarded permanently (Stich et al. 2018; Lin et al. 2018). Wire-byte savings depend on the encoding, while convergence depends on the residual handling, optimizer, and compression schedule; a value-count reduction alone establishes neither.

↩ Gradient quantization, top-k sparsification, and error feedback that carries unsent updates forward are established in section 5.9. Convergence depends on the optimizer and compression assumptions.

Selective update sharing further reduces communication by transmitting only subsets of model parameters or updates. In layer-wise selective sharing, clients update only certain layers, typically the final classifier or adapter modules, while keeping the majority of the backbone frozen. This reduces the update payload and limits which shared representations can change, but the trainable subset can still overfit nonrepresentative client data.

Architectural partitioning requires specifying what crosses the boundary. A personal/shared parameter partition keeps local modules private while aggregating selected shared parameters. Split learning instead divides execution between client and server, exchanging intermediate activations and backward gradients at the cut. These designs have different traffic and exposure patterns; neither guarantees a communication reduction or privacy merely by dividing the model.

These approaches can use the FedAvg aggregation protocol (section 11.6.2.2) only when their encoding and shared parameter schema support weighted aggregation. Under secure aggregation, which sums masked updates without inspecting them, weighted aggregation also requires every client to share one quantization scale and one coordinate template for sparse values (section 11.7.4). Local-only parameters must stay outside the shared aggregate altogether.

Whatever the encoding, the total communication volume for round r is

$$M_r = M_{\text{download},r} + M_{\text{upload},r} + M_{\text{protocol},r} + M_{\text{retry},r}, \tag{11.4}$$

Equation 11.4 makes download, upload, protocol, and retry traffic explicit. The download distributes model state, the upload carries client updates, protocol traffic covers authentication and cryptographic setup, and retries account for failed or repeated transfers. Optimizing only $M_{\text{upload},r}$ can therefore move rather than remove the bottleneck. Additional local computation amortizes network transfers, and each communication-efficient alternative to it carries its own cost. Compression may degrade gradient fidelity, selective updates can limit model capacity, and split architectures complicate coordination. Summed over a training run, the upload term alone can exceed the raw data that federation was meant to keep local, even when each update is far smaller than that data.

↪ The break-even condition under which encoding and decoding time costs less than the transfer time compression saves is developed, with a worked example, in section D.5.

Napkin Math 11.3: Model updates vs. raw data

Problem: A team is designing a federated camera-personalization system that learns from 200 MB of compressed user images per week. Compare one upload of those images with one compressed model-update uplink round.

Math: Bandwidth efficiency is the ratio of raw data volume to model update size.

1. **Raw data upload**: 200 MB/week.
2. **One federated uplink round**: A compressed model update (5M params) is 2.5 MB.
3. **One-round uplink reduction**: 200 MB/2.5 MB = 80×.

Systems insight: For this scenario, one 2.5 MB uplink is 80× smaller than one 200 MB raw-data upload. That ratio is not the end-to-end communication saving of a training run: after roughly 80× uplink rounds, cumulative update uploads alone equal the raw-data volume, before counting global-model downloads, retries, or protocol overhead. The verdict therefore turns on the round schedule. A deployment that runs one round per day uploads seven updates per week of data, far below the break-even count, and saves most of the uplink traffic. One that runs a round every hour uploads 168 updates per week, more than the break-even count, and moves more bytes than shipping the images would. Per-round traffic must be multiplied by the expected round count before federation can claim a bandwidth saving.

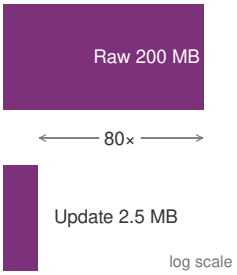

One compressed update is 80× smaller than a week's raw-image upload.

The local-work schedule and update encoding jointly determine communication cost. Figure 11.14 compares uncompressed 32-bit values (Panel A) with uniform 8-bit quantization (Panel B), which gives a $4\times$ reduction in value bytes before scales and other metadata are counted. Top-k sparsification with retained error feedback (Panel C) must encode coordinate identities as well as selected values. Error feedback retains one residual value per update coordinate between transmissions. Its storage depends on the update size and residual precision and must fit the device allowance alongside training and replay state. The receiving path must reconstruct or aggregate these encodings

consistently, with separate compatibility requirements when secure aggregation hides individual updates.

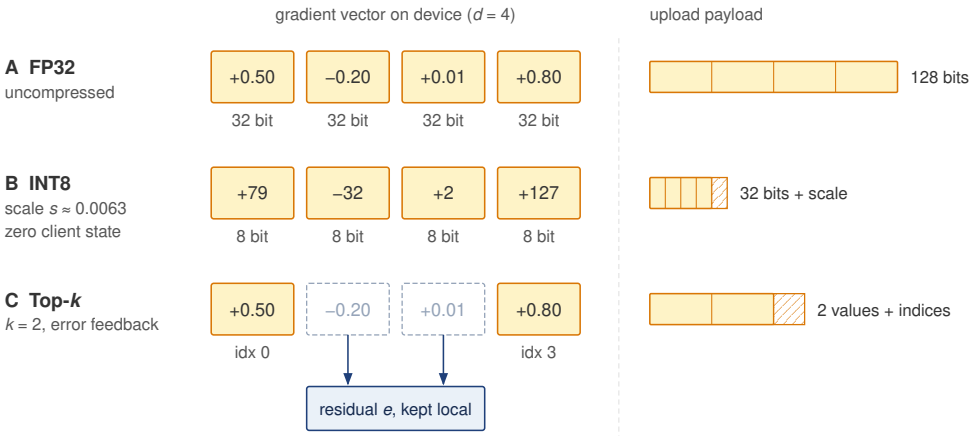

Figure 11.14: Gradient Compression Techniques: Standard updates transmit full-precision values; quantization maps values to lower-precision buckets, such as FP32 to INT8; and sparsification transmits only the most significant top-k gradients. The latter two mechanisms reduce bytes transferred but introduce representation error or residual state that the training algorithm must manage.

The quantization panel (Panel B) is labeled "zero client state" because it keeps no persistent error-feedback residual; encoding buffers, scales, and temporary workspace still enter the device budget. Compression is useful when its reduction in total communication outweighs encoding work and any extra rounds needed to reach the quality target. A smaller upload helps a round only if transfer was what made clients miss the deadline or excluded a cohort, and the round's completion logs show which. Smaller payloads can let clients on slow or congested uplinks finish before their resource window closes, so compression is judged by completed participation and cohort outcomes, not by bytes saved alone. Even when transfer becomes feasible, the shared model may still fail to capture local variation.

11.7.3 Federated personalization

Compressing updates lets more devices finish a round, but clients that complete compatible updates can still pull the shared model in opposite directions when a dialect phrase, a regional keyboard shortcut, or an individual's wrist-worn sensor placement conflicts with the population mean. A single global model may serve the average user adequately while concealing persistent accuracy losses for specific cohorts. Personalized federated learning retains shared representation learning while allowing selected parameters or models to specialize locally, trading added local state, training cost, and overfitting risk on sparse user data against per-user quality.

Let θ_k denote the model parameters on client k, and θ_{global} the aggregated global model. For N_{clients} clients, counted as in equation 11.3, traditional FL seeks to minimize a global objective:

$$\min_{\theta} \sum_{k=1}^{N_{\text{clients}}} w_k \mathcal{L}_k(\theta)$$

where $\mathcal{L}_k(\theta)$ is the local loss on client k, and w_k is a weighting factor (e.g., proportional to local dataset size). This formulation optimizes the weighted average loss of a shared model θ; it does not guarantee acceptable performance for every client. Local loss landscapes $\mathcal{L}_k$ can differ across clients, reflecting non-IID data distributions and varying task requirements.

Personalization modifies this objective to allow each client to maintain its own adapted parameters θ_k, optimized with respect to both the global model and local data:

$$\min_{\theta_1,\ldots,\theta_{N_{\text{clients}}}} \sum_{k=1}^{N_{\text{clients}}} \left(\mathcal{L}_k(\theta_k) + \lambda_{\text{reg}} \cdot \mathcal{R}(\theta_k, \theta_{\text{global}}) \right)$$

Here, $\mathcal{R}$ penalizes deviation from the global model, for example with the squared distance that FedProx adds to each local loss, and λ_{reg} controls the strength of that penalty. The per-client parameters θ_k are thus state that the plain federated objective does not have, held near the shared model by the penalty rather than trained on local data alone. The displayed objective describes local specialization around a given global reference; a complete protocol must additionally specify how that reference is trained and refreshed. The penalty trades local fit against proximity to the shared model, without guaranteeing improvement for every client.

In a wearable health monitor that classifies physical activities from physiological signals, this objective separates shared motion features from individual gait signatures, resting baselines, and sensor placements. Finetuning the final classification layer or low-rank adapters around the global reference θ_{global} addresses those user-specific differences, provided evaluation establishes that the local examples actually improve accuracy on rare activity classes with little supervision.

Several personalization strategies address the trade-offs between compute overhead, information exposure, and adaptation speed. One widely used approach is local finetuning, in which each client downloads the latest global model and performs a small number of gradient steps using its local data. This method avoids transmitting those examples during adaptation, but the resulting local state and any later update remain privacy-sensitive. It may also yield suboptimal results when the global model is poorly aligned with the client's data distribution or when the local dataset is extremely limited.

Another effective technique involves personalization layers, where the model is partitioned into a shared backbone and a lightweight, client-specific head, typically the final classification layer (Arivazhagan et al. 2019). One such method, FedPer, has clients train both the shared base and their personal head, but only the shared base participates in server aggregation; the head remains local. A separate head-only adaptation policy can freeze a downloaded backbone to reduce local training cost when the relevant variation lies mainly in output categories or decision boundaries.

Clustered federated learning offers an alternative by grouping clients according to similarities in their data or performance characteristics, and training separate models for each cluster. This strategy can enhance accuracy within homogeneous subpopulations, though it introduces additional system complexity and may require exchanging metadata to determine group membership.

Finally, meta-learning approaches, such as model-agnostic meta-learning (MAML),[15] aim to produce a global model initialization that can be quickly adapted to new tasks with a few local updates (Finn et al. 2017). A meta-learned initialization is especially useful when clients have limited data or operate in environments with frequent distributional shifts.

Each strategy occupies a different point in the personalization trade-off space. Local finetuning, personalization layers, clustered federated learning, and meta-learning differ in compute overhead, information exposure, and adaptation latency (table 11.7). The mechanism determines what state is shared; a privacy guarantee requires separately specified controls and assumptions.

Table 11.7: Personalization Trade-Offs: Federated learning strategies balance personalization with compute, information exposure, and adaptation latency. Data locality narrows exposure but does not by itself provide a privacy guarantee.

Strategy	Personalization mechanism	Compute overhead	Privacy preservation	Adaptation speed
Local Finetuning	Gradient descent on local loss postaggregation	Low to Moderate	Raw examples stay local; adapted state remains sensitive	Fast (few steps)
Personalization Layers	Split model: shared base + user-specific head	Moderate	Narrows shared state; does not guarantee privacy	Fast (train small head)
Clustered FL	Group clients by data similarity, train per group	Moderate to High	Group membership and updates can leak attributes	Medium

[15] **MAML (model-agnostic meta-learning)**: Finn et al. (2017) optimize for a model initialization from which one or a few gradient steps can yield good performance on a new task, with experiments including one-shot and few-shot settings. Meta-training is more expensive than ordinary finetuning because it differentiates through the adaptation process, but deployment adaptation can be limited to a small number of local updates. This makes MAML relevant for edge scenarios where inference compute is cheap but training data is scarce.

Strategy	Personalization mechanism	Compute overhead	Privacy preservation	Adaptation speed
Meta-Learning	Train for fast adaptation across tasks/devices	High (meta-objective)	Depends on meta-training and adaptation protocol	Very Fast (few-shot)

Selecting an on-device personalization strategy requires balancing adaptation capacity against memory residency. Three fine-tuning paradigms differ in structural footprint (figure 11.15). Full fine-tuning updates every backbone layer (Column A), head-only adaptation freezes the feature extractor (Column B), and parameter-efficient adapters sit beside frozen backbone blocks (Column C).

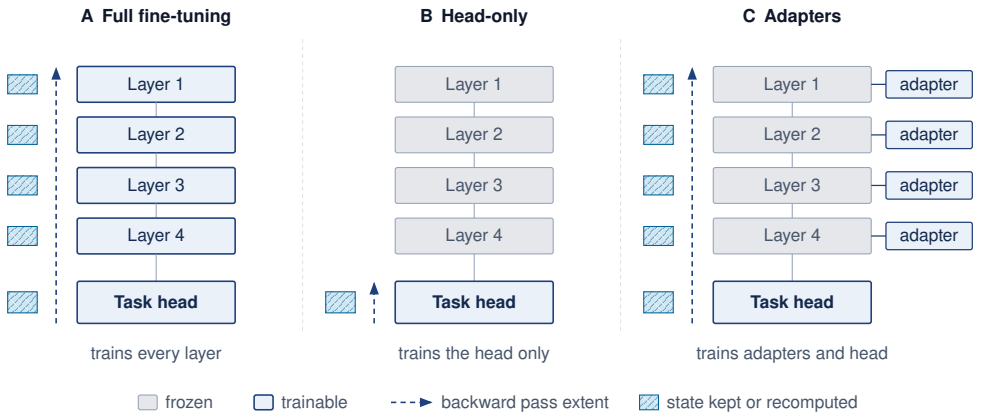

Figure 11.15: Federated Personalization Architectures: (A) Full fine-tuning updates every layer; (B) head-only adaptation freezes the feature extractor and trains only the task head; (C) adapter-based learning, such as LoRA, inserts small trainable modules beside a frozen backbone. The dashed arrow marks how far the backward pass must reach. Adapters shrink trainable state, but because they sit at every depth their gradients still traverse the backbone, so the required backward state must be retained or recomputed. In this comparison, head-only adaptation truncates the backward pass at the feature extractor.

The dashed backward-pass arrows in figure 11.15 are the activation term of equation 11.1 drawn as distance. Adapters at early layers in Column C pull the backward path through the later backbone, whose saved state must be retained or recomputed, while head-only adaptation in Column B ends it at the head, as in the wake-word case of section 11.5.1. Placement sets the dependency path, and the execution schedule sets how much of that state stays resident. Keeping the personal head or adapter on the device also narrows what crosses the device boundary, but the shared backbone updates that still leave the client remain derived from private local examples and require explicit privacy protection.

11.7.4 Federated privacy

Even when personal parameters and raw training examples never leave the phone, an adversary or curious server that inspects an uploaded gradient or weight delta $\Delta\theta$ can reconstruct typed text, recover facial features, or determine whether a specific phrase appeared in a user's local dataset. Model inversion attacks[16] and membership inference attacks exploit the fact that transmitted updates encode the very examples that produced them.

Two protective measures address that leakage. **Secure aggregation**[17] protocols reveal a combined result while hiding individual updates beyond what that result itself reveals, subject to the protocol's assumptions about collusion and participating clients. Differential privacy, a formal bound on how much any one protected user's or record's inclusion can change a released result, is obtained by clipping each contribution and adding noise calibrated to that bound. The deployment must specify which unit receives that protection.

These controls solve different problems. Secure aggregation limits what the coordinator learns about each submitted update, whereas differential privacy bounds how much a declared privacy

16 **Model inversion attack**: Reconstructs training data from model outputs or gradients by optimizing inputs that maximize confidence scores (Fredrikson et al. 2015; Zhu et al. 2019). In federated learning, gradient updates can leak more information than model outputs alone because they expose training-step information directly. Defenses such as gradient perturbation, secure aggregation, and differential privacy add protocol, utility, communication, or compute trade-offs whose magnitude depends on the threat model and deployment; section 13.8.5.2 develops the threat model for these attacks.

17 **Secure aggregation**: Cryptographic protocol where client pairs exchange random masks that cancel in the server-side sum, revealing only the aggregate gradient without exposing individual contributions (Bonawitz et al. 2017). Production-scale federated systems such as Gboard combine secure aggregation with eligibility checks, dropout handling, and round scheduling (Bonawitz et al. 2019). The communication expansion and minimum-participation requirements depend on protocol parameters and the security model, so aggregation systems must wait for enough eligible devices before releasing a round.

unit can affect a released result. A deployment therefore needs a privacy accounting log that records the unit of protection, clipping rule, noise mechanism, sampling assumptions, (ϵ, δ) expenditure for every release, and cumulative composition across training and federated analytics (Abadi et al. 2016; Kairouz and McMahan 2021). Clipping bounds each contribution before noise is calibrated. The parameters ϵ and δ describe the permitted change in output probabilities when the protected user or record is included or removed; composition accounts for repeated releases involving that unit. Resetting a round counter does not reset privacy loss for a user who participates again.

↪ The composition rules that turn a privacy accounting log's per-release entries into one cumulative privacy budget are developed in section 13.9.2.

Figure 11.16 models this cryptographic boundary through pairwise zero-sum masking. Clients A and B negotiate a shared random mask $\mathbf{s}$ via Diffie-Hellman key agreement, which lets the two clients derive the same secret from public values that the server relays but cannot use to compute it. They perturb their local updates into pseudo-random vectors ($\mathbf{y}_A = \Delta\theta_A + \mathbf{s}$ and $\mathbf{y}_B = \Delta\theta_B - \mathbf{s}$) before crossing into the server domain. When the aggregation server sums the received vectors, the symmetric masks cancel ($\mathbf{s} - \mathbf{s} = 0$), yielding the exact unmasked aggregate $\Delta\theta_A + \Delta\theta_B$. This two-client illustration assumes both clients complete the exchange and neither reveals its mask to the server. It demonstrates cancellation, not a complete protocol for collusion resistance or dropout recovery.

A real round has many clients, and some drop out after agreeing on masks but before uploading. Each dropped client leaves its pairwise masks uncancelled in the sum. The full protocol therefore has every client split the secret seeds of its masks into shares held by the other clients, and any threshold number of those shares can rebuild a seed (Bonawitz et al. 2017). After the upload deadline, surviving clients send the shares of the dropped clients' seeds, and the server regenerates and subtracts the dangling masks. Recovery therefore costs an extra message round and computation that grows with the number of dropouts, and it succeeds only if at least the threshold number of clients survive; otherwise the round is abandoned. Over-selection adds to this cost. A client that completed key agreement but whose report the coordinator discards once the target count arrives is a dropout to the protocol, so every discarded report adds recovery work, and the clients kept must still clear the threshold.

That threshold is one reason for a participation floor. The other is the sum itself, because an aggregate over a few clients reveals too much about each of them. The coordinator must therefore invite enough clients that the expected survivors clear the floor, and a cohort that falls short is held for a later round rather than aggregated.

Masking also constrains how clients encode their updates (section 11.7.2). With cleartext uploads, the server can dequantize each client's tensor with that client's scale factor and scatter sparse index-value pairs into a dense buffer before summing. Under secure aggregation the masks are added in modular integer arithmetic, so each upload is a fixed-point integer vector, and the server sums these masked vectors element by element without inspecting individual uploads. Clients that quantize updates must therefore quantize to a shared fixed-point scale so their masked integers represent the same unit at every coordinate, and clients that sparsify updates must project onto a shared coordinate template or sketch rather than uploading client-specific index lists that would fail to align and would reveal which parameters changed.

Secure aggregation also creates a robustness conflict. Byzantine-resilient rules such as Krum, coordinate-wise filtering, or per-client anomaly inspection need access to individual updates, while secure aggregation deliberately hides those updates (Blanchard et al. 2017; Bonawitz et al. 2017). A system cannot claim both arbitrary per-client inspection and conventional secure aggregation over the same messages. It must instead validate clients before masking, use attested execution and bounded local updates, compute only aggregation-compatible checks, or adopt a protocol designed to evaluate approved robust statistics without revealing raw client contributions.

↩ Byzantine failures, in which a worker stays alive but returns wrong results, and aggregation rules such as Krum and trimmed mean that bound their influence are established in section 7.3.3.

Secure aggregation and differential privacy can protect different boundaries, but their composition must follow the threat model. If a trusted server adds noise after receiving an exact aggregate, the privacy guarantee applies to the released noisy result; it does not hide that exact aggregate from the server. A protocol designed for distributed noise can instead let the server receive only a noisy aggregate. Clipping, noise calibration, dropout assumptions, and repeated-release accounting must match the chosen protocol. Neither mechanism establishes that local data or updates are trustworthy.

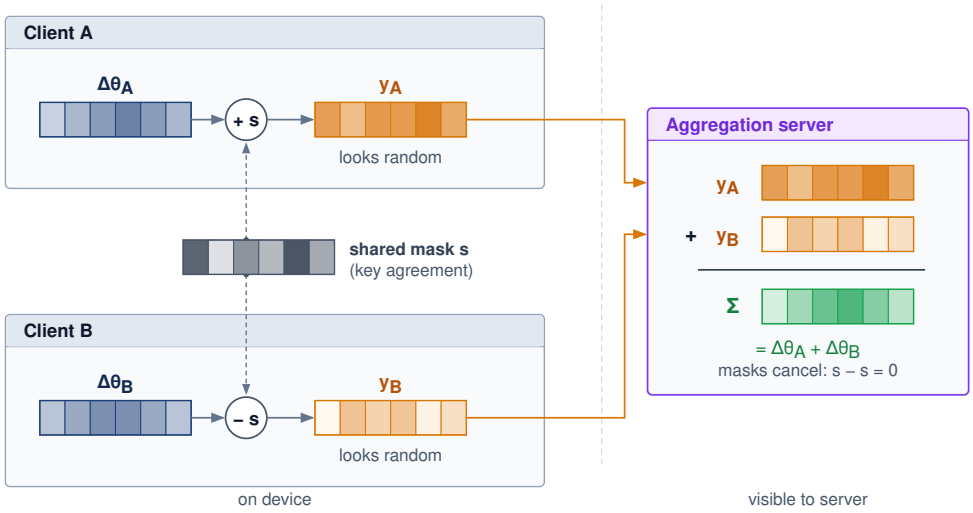

Figure 11.16: Secure Aggregation Protocol: A simplified two-client view of pairwise masking. Clients A and B agree on a shared random mask $\mathbf{s}$; A uploads $\mathbf{y}_A = \Delta\theta_A + \mathbf{s}$ and B uploads $\mathbf{y}_B = \Delta\theta_B - \mathbf{s}$, so each upload alone looks random to the server. Summing them cancels the mask and yields exactly $\Delta\theta_A + \Delta\theta_B$. The server learns only the aggregate, which can still leak information unless differential privacy bounds it.

> **Checkpoint 11.3: Federated system design**
>
> Consider architecting a federated learning system for a fleet of 10 million mobile devices. The data is highly non-IID (users have distinct, clustered typing patterns), and the network environment is constrained (1–10 Mb/s). Before choosing optimizations for this scenario, classify the local bottlenecks and state which system constraint dominates each decision.
>
> Consider three design decisions:
>
> 1. Aggregation: Select plain FedAvg, reduced local epochs, or personalization/clustering for the non-IID population, and justify how the choice controls client drift.
> 2. Privacy: Decide whether the deployment needs formal differential-privacy noise or whether secure aggregation and data minimization are sufficient for the stated threat model. Using the clipping and noise described in this section, explain how noise can affect utility and rounds to the same quality target, and what measurements would establish that cost.
> 3. Communication: Choose between gradient quantization (e.g., 8-bit) and structured sparsity under constrained bandwidth, and justify the bandwidth/accuracy trade-off.

11.7.5 Large-scale device orchestration

A round that selects a million phones overnight will see some fraction overheat, lose connectivity, or be unplugged before they report, and a few may submit corrupted or deliberately malicious updates. Coordinating that population requires protocols that handle Byzantine clients, network partitions that temporarily cut participants off, and constrained communication. These challenges differ from controlled data center training, where high-bandwidth networks and managed infrastructure narrow the trust and availability envelope.

11.7.5.1 Asynchronous device synchronization

Federated learning inherits failures from distributed systems and machine learning, but its coordination objective is not consensus, the problem of making replicated servers agree on one ordered log of operations. Federated clients agree on no such log. A coordinator instead decides which

stochastic client updates are eligible, how stale updates are weighted, and which aggregation rule can tolerate faulty participants without letting them steer the model. Devices crash, lose power, or disconnect during rounds, while compromised clients can submit poisoned updates crafted to push the shared model toward an attacker's goal. Robust aggregation can reduce the influence of some faulty participants, but its assumptions, compute cost, and compatibility with secure aggregation must be explicit.

The trust chain begins before aggregation. A production coordinator authenticates the device and application instance, verifies the signed model and training-policy version, checks eligibility and freshness, binds the returned update to that version and round nonce, and records whether the update passed protocol-level bounds. Hardware-backed attestation can strengthen claims about code identity on supported devices, but it cannot prove that local examples are benign or representative. The resulting update-integrity record supports rejection, rollback, and incident analysis without treating device identity as evidence that the gradient is correct.

Network partitions are the harder case. Unlike a data center network, a federated fleet must absorb clients that stay offline for hours or days while traveling, out of coverage, or powered down. Asynchronous coordination protocols keep training moving despite missing participants, but they must balance staleness (accepting outdated contributions) against freshness (prioritizing recent but possibly sparse updates). Buffered asynchronous aggregation (Nguyen et al. 2021) shows the mechanism. It removes the round barrier, letting each client train whenever its admission test passes and upload when done, while the server accumulates arriving updates in a buffer and applies their aggregate once a fixed number have arrived. Each update is down-weighted by its staleness, the number of server versions published since the client downloaded its starting model, so a slow client still contributes without dragging the shared model back toward an old state.

Conventional pairwise secure aggregation (section 11.7.4) does not fit this buffer. Its masks cancel only among the clients fixed at key agreement, whereas the buffer fills with whichever clients finish. Staleness weighting breaks the masks as well. Weighting a masked update on the server also scales its masks, so they no longer cancel, and the client cannot apply the weight before masking because its staleness is known only when the update arrives. A buffered design therefore needs a secure-aggregation protocol built for it.

Recovery must distinguish coordinator state from device-local state. The coordinator checkpoints the global model and the optimizer and protocol state needed to resume accepted progress; this does not require a simultaneous checkpoint on every device. An incomplete round needs an explicit acceptance or abandonment rule, followed by evaluation of the population represented by its completed updates. State reconciliation identifies a returning device's base version and determines which shared state it must download and which personal state remains compatible. Server-side load balancing can redistribute coordination work, but replacing an unavailable client with a faster one also changes the local data contributing to training. Whereas convergence guarantees for synchronous training assume that every participant completes each round, a federated coordinator must treat stragglers[18] and dropouts as the normal case and state what its convergence argument assumes about them.

↩ Gradient staleness, the delay between the model version a worker computed against and the version its update modifies, and its effect on convergence under bounded and fully asynchronous synchronization are established in section 6.4.4.

[18] **Straggler problem**: In data center training, stragglers slow synchronous progress; in federated learning, the problem is amplified by heterogeneous hardware, intermittent connectivity, and eligibility windows (Bonawitz et al. 2019; Kairouz and McMahan 2021). Solutions like asynchronous aggregation and bounded staleness (ignoring updates computed against a model more than a fixed number of rounds old) improve throughput but introduce a fairness trade-off: aggressive straggler mitigation biases models toward users with high-end devices, potentially degrading performance for the most constrained users who need personalization most.

11.7.5.2 Heterogeneity management at fleet scale

Million-device orchestration is a tiering decision, in which the system decides which clients can train, which can only report lightweight updates, and how much bias that selection introduces. The difficulty is that hardware capabilities, network conditions, data distributions, and availability patterns vary simultaneously, challenging traditional distributed machine learning assumptions about homogeneous participants operating under similar conditions.

Illustrative peak compute ratings differ by about 1,200× between flagship smartphones rated at 35 TOPS and, below Archetype C's phone-class clients, IoT microcontrollers rated at 0.03 TOPS. These ratings do not translate into a training-speed ratio, because supported precision, backward operators, memory traffic, and sustained execution determine update time. Memory spans the same 65,536× range from microcontroller SRAM to flagship phone RAM met in section 11.3.2, and at fleet scale that range decides which update state each tier can retain at all. Energy budgets further restrict when that work can run. Background adaptation in this example targets roughly 500–1000 mW, while phones may sustain 2–3 W for foreground ML workloads and burst higher briefly. Network

throughput, latency, and interruption patterns must likewise be measured for the participating clients before selecting update frequency and compression.

Measured scores show a spread of the same order among phones alone, without counting microcontrollers. The AI Benchmark from ETH Zurich runs a fixed suite of quantized and floating-point inference workloads on each phone processor through its Android acceleration path (Ignatov et al. 2019; ETH Zurich Computer Vision Lab 2026). Among processors released from 2020 on, which a client fleet can still contain, the scores span about 400×, from the Snapdragon 460 (2020) to the Dimensity 9600 Pro (2026), and figure 11.17 plots every processor in the ranking. Even flagships of the same generation differ by 14×, with the Google Tensor G4 scoring that far below the Dimensity 9600 M, largely because each score reflects the processor's software path as well as its silicon. The Tensor G4 result runs through Android's generic neural-network interface, while the leaders use their vendors' own runtimes. These are inference scores, not training speeds, but they bound the problem from the side that matters for scheduling. No single local-work quota fits devices that differ by two orders of magnitude in measured throughput.

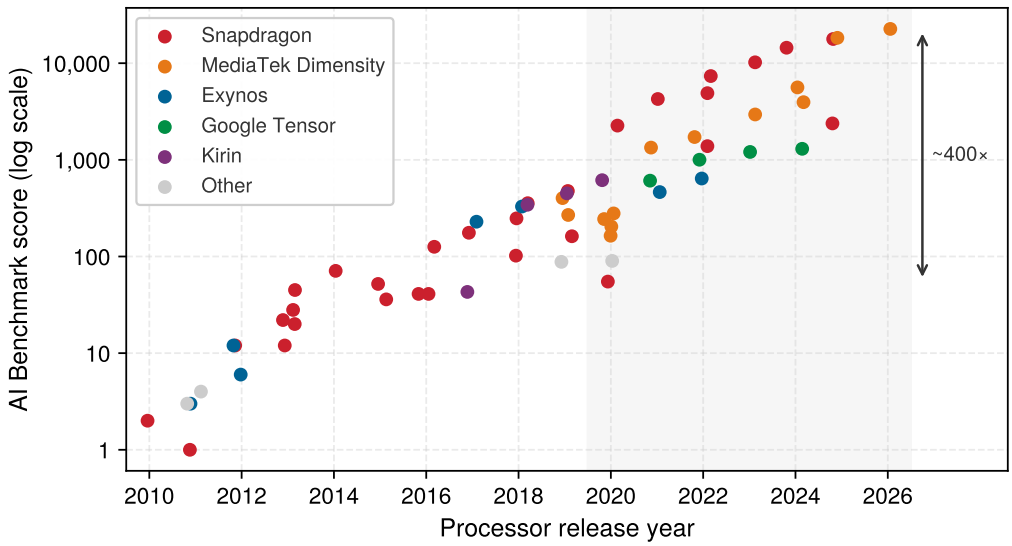

Figure 11.17: Measured AI Performance across Phone Processors: Aggregate AI Benchmark score of mobile systems-on-chip by release year on a logarithmic axis, colored by processor family. The shaded window marks processors from 2020 onward that a client fleet can still contain; the bracket gives the spread of scores within it. Scores measure inference through each processor's software acceleration path, not training speed. Data from ETH Zurich Computer Vision Lab (2026), method of Ignatov et al. (2019).

Tiered participation can assign different amounts of local work while preserving a compatible shared-update contract. A capable phone may complete more steps, while a constrained client contributes a smaller update or a separately defined aggregate statistic. These outputs are not interchangeable simply because they come from the same fleet. More work on strong devices also cannot replace the observations held by weak devices, so the policy must evaluate both computational contribution and population coverage.

The scale challenges extend beyond device heterogeneity to coordination overhead. Federated learning does not run a consensus protocol among millions of clients, nor does hierarchical aggregation perform local consensus. Hierarchical aggregation reduces fan-in through regional aggregators, which validate an eligible cohort and combine its updates before forwarding a partial aggregate; agreement on an ordered replicated log, if needed for the coordinator's own metadata, remains a separate control-plane service. Edge infrastructure can provide aggregation points, but each added tier creates another trust boundary and another place where weighting, clipping, privacy accounting, and failure recovery must remain consistent.

Tiering interacts with the participation bias analyzed under client scheduling (section 11.7.1). Capability-based selection shortens rounds but concentrates completed updates on strong devices, so

stratified sampling across tiers must be paired with per-cohort evaluation of completed participation and model quality. Availability also has a daily rhythm; devices that charge overnight in one time zone form a different cohort from those available during another zone's working hours.

Orchestration brings client scheduling, update compression, personalization, privacy, and tiering into a coherent policy for a single federated round. Deployed systems, however, do not stop after one round, and their devices keep applying local and federated updates week after week.

11.8 Continual Adaptation Under Edge Budgets

A keyboard model that adapts each night for a year passes through hundreds of admitted update windows, and each window moves it a little further from the version that was validated at release. Every update the chapter has built so far was judged one at a time, by whether it fit the memory budget, passed the admission test, and improved a local metric. Repeated across a deployment's lifetime, those individually acceptable updates compound. Earlier changes alter what later updates start from, the replay buffer decides which old observations still count, and participation decides which devices shape the next shared version.

The mechanisms that keep this trajectory stable are the ones already introduced, now applied over time. Selective updates bound how far any one window can move the model, and they save work only when the execution path can skip it. Replay protects earlier behavior at the cost of retained examples and additional training. Opportunistic admission delays learning until device conditions permit it. Each needs a resource budget and an outcome check on retained behavior alongside the newest data. The long-lived stream that drives this learning is mostly unlabeled, which sets the first design question.

11.8.1 Learning from unlabeled sensor streams

A camera stream running continuously at 30 frames per second produces roughly 2.6 million frames daily before retention filtering. An accelerometer sampling at 100 Hz generates 8.6 million readings, almost none of them accompanied by a user label. Cameras, accelerometers, location sensors, and touch interactions all supply data that can support self-supervised learning, turning continual learning into a storage and scheduling decision. The engineering question is which traces encode stable structure worth learning from and which should expire before they crowd out model state, replay buffers, or user-facing storage.

Learning from such traces without labels requires the supervision to come from the data's own structure. Contrastive learning is the common way to manufacture it on a sensor stream, training an encoder so that the embeddings of designated positive pairs, two samples that should mean the same thing, lie close together while those of designated negative pairs lie far apart. Nearby camera frames or augmented views of one frame can supply positives when they preserve the task-relevant content; scene changes can invalidate that assumption. Different frames or replay entries can supply negatives, but may depict the same underlying event. The pair-selection policy therefore defines a learning signal that needs validation. Other devices' samples are not automatically available under a federated data-locality policy, and stored pairs or representations compete with replay and model state for memory.

11.8.2 Continual adaptation and catastrophic forgetting

A model that adapts continually to recent summer lighting, a new work commute, or a shift in user vocabulary can overwrite the weights it needed for winter weather or older commands if each update moves parameters purely to minimize recent loss. When a deployment uses continual adaptation to address changing environments, user behavior, or task requirements, each update must balance new evidence against previously learned behavior. This is the stability-plasticity trade-off: models must remain stable enough to preserve useful knowledge while plastic enough to learn new patterns.

Continual learning on edge devices faces several interconnected challenges. Catastrophic forgetting (section 11.5.2) is magnified when devices cannot store or revisit historical training data. Task interference compounds it when several learning objectives compete for limited parameter capacity, and distribution shift between pretraining and local deployment demands continued adaptation without eroding the model's core representations. Restricted device memory rules out the comprehensive replay buffers that centralized infrastructure can afford.

Three families of mechanisms manage the stability-plasticity trade-off during continual learning. Elastic weight consolidation[19] (Kirkpatrick et al. 2017) mitigates catastrophic forgetting by penalizing changes to weights important for previous tasks, where a weight's importance reflects how sharply the earlier tasks' loss rises when that weight moves. Experience replay (section 11.5.2) maintains stability during adaptation by interleaving new training with replayed examples from previous tasks. Progressive neural architectures allocate new capacity for new tasks rather than forcing all knowledge into fixed-capacity networks (Rusu et al. 2016). That last option is attractive when interference is severe, but it consumes additional memory and therefore fits only devices with enough spare capacity.

[19] **EWC (elastic weight consolidation)**: Estimates each weight's importance with a diagonal approximation to the Fisher information, a measure of how sensitive the earlier tasks' loss is to that weight, and uses those estimates to weight a penalty on changes from previously learned parameters (Kirkpatrick et al. 2017). In addition to current parameters, a basic implementation stores importance scores and reference parameter values. Retaining separate penalties for earlier tasks can increase that state further. Compared with replay, EWC trades stored examples for parameter-sized statistics; which uses less memory depends on model size, retained history, precision, and the penalty implementation.

The earlier mechanisms bound each step without guaranteeing the trajectory. Meta-learned initializations (section 11.7.3) and few-shot adaptation (section 11.5.1) reach useful local behavior in few updates. Restricting updates to bias terms, adapters, or heads (section 11.4) keeps the base parameters intact, while useful pretrained representations reduce what the device must learn (Bommasani et al. 2021). Even a single update to a head or adapter can alter predictions on earlier tasks, and a year of individually small updates can accumulate into significant behavioral drift. The acceptance test for each window must therefore measure retained capability alongside improvement on the new data, against a recoverable baseline, because neither a small update count nor a frozen backbone is a stability guarantee.

Once the update footprint is bounded, the runtime must decide when the update can execute and whether it should be shared. The admission condition in equation 11.2 already governs when a continual update may run; what changes over a long deployment is that each admitted window also moves the device further from its last validated state. When the device also shares updates, the compression and selective transmission of section 11.7.2 keep that repeated traffic inside the bandwidth and energy budgets (Konečný et al. 2016).

Stability then becomes a state-management problem. Replay buffers, support sets, adaptation logs, and model update metadata must be protected against unauthorized access or tampering because they encode the local evidence that shaped the model. Lightweight encryption or hardware-backed secure storage can reduce that exposure, but security controls do not prove that adaptation remains beneficial. Lightweight validation techniques, including confidence scoring, drift detection heuristics (Gama et al. 2014), and shadow model evaluation, monitor adaptation dynamics and can hold an update for validated review before severe degradation occurs. Robust rollback procedures require a trusted baseline checkpoint that can be restored quickly, especially in safety-important and regulated domains where failure recovery must be provable.

Privacy and compliance requirements close the loop. User consent, data minimization, retention limits, and the right to erasure must be designed into the adaptation pipeline rather than added after deployment. Meeting regulatory obligations at scale demands on-device learning workflows in which the model can adapt in place while the operator can still explain, bound, and reverse each change. Because a fleet applies these rules on flagship phones, budget phones, and microcontroller sensors at once, the pillars and their stability controls must be assembled into one policy whose terms change with the device that runs it.

11.9 Integrating the Edge Intelligence Pillars

Consider the voice assistant from the opening of this chapter, now deployed across 50 million devices that range from flagship phones to budget handsets. The chapter built model adaptation, data efficiency, and federated coordination as separate pillars, and treated separately, each solves its own constraint while leaving the others free to misbehave. An adapter rank chosen without regard to replay memory can crowd out the buffer that protects earlier behavior, and an update format chosen without regard to aggregation can produce arrays no coordinator can combine. The integration that works makes one per-tier policy decide how much a device may adapt, how much it may remember, and how it may participate in the fleet, because all three draw on the same memory, energy, and connectivity budget.

The model adaptation layer starts with candidate policies for the scenario's device tiers. Flagship phones (Tier 1) representing 20 percent of the deployment receive LoRA adapters of rank 32. Mid-tier devices (Tier 2), comprising 60 percent, receive rank 16 adapters. Budget devices, the remaining 20 percent, use bias-only updates under a 1 GB memory limit. These assignments apply the device-tier

framework of section 11.4.3.4. Budget phones receive the Tier 3 mechanism because the memory granted to background adaptation, not installed RAM, sets their tier. The voice assistant keeps one adaptation per user on a latency-sensitive speech path, so it uses mergeable LoRA on both capable tiers rather than the switchable residual adapters that suited the multi-context workload. These assignments specify an initial policy, not a demonstrated fit. Profiling must include the base model, activations, optimizer state, runtime, and foreground reserve, while local validation determines whether each update space captures useful voice variation.

The data efficiency layer allocates part of the same device budget to experience replay, with candidate buffers of 10 MB on budget devices and 100 MB on flagship models. Buffer coverage and retention tests determine whether these allocations preserve earlier behavior. The policy targets initial adaptation within 5–10 interactions, addressing the cold-start problem of limited evidence for a new user. That target requires enough valid feedback, since a handful of ambiguous interactions supervises nothing. Later streaming updates are admitted only when both learning evidence and device resources permit them.

The federated coordination layer combines explicitly compatible shared state. Different adapter ranks and bias-only policies do not produce interchangeable arrays, so the coordinator needs separate compatible groups or a defined mapping to common parameters. Even equal-rank LoRA factors require a stated aggregation rule because each client's effective weight update is the product of its up-projection and down-projection matrices. Averaging the two factor matrices separately across clients and then multiplying the averages does not give the average update. The product of the averages contains cross-terms that pair one client's up-projection with another client's down-projection, combinations that no device ever trained. Averaging the multiplied weight updates instead preserves the exact mean update, but a sum of several low-rank products generally has higher rank, so it must be folded into the base weights or truncated back to the target rank before the next round. Tier-specific personal state can remain local. The scenario's 50 MB adapter payload is smaller than the 14 GB full-model payload, but downloads, repeated rounds, and protocol overhead still enter the communication budget. Each participating device must pass the energy, thermal, and foreground-latency admission tests. Raw audio can remain local while derived updates and participation metadata remain sensitive. Secure aggregation, privacy accounting, access controls, and retention policies must cover that shared state.

These layers only work as a system when the integration policy protects against capability mismatch, component failure, resource conflicts, and untested emergent behavior. Four edge-learning integration controls answer those risks one for one in table 11.8: hierarchical capability matching, graceful degradation, explicit priority policy, and performance validation.

Table 11.8: Edge-Learning Integration Controls: A tiered edge-learning deployment needs explicit controls for capability matching, graceful degradation, resource conflicts, and validation across device and network combinations.

Integration risk	Control	Why it matters
Capability mismatch	Hierarchical capability matching	Deploys more sophisticated techniques on capable devices while preserving basic functionality across the full device spectrum.
Component failure	Graceful degradation	Keeps local adaptation useful when connectivity is poor or battery constraints force the system into a minimal adaptation mode.
Resource conflict	Explicit priority policy	Prevents model adaptation and replay buffers from competing for the same memory, energy, or latency budget without a predefined owner.
Emergent behavior	Performance validation	Tests combinations of device tier, network state, and adaptation policy because integrated systems fail in ways individual techniques do not expose.

The same reasoning reduces to three questions (figure 11.18). The first asks how much adaptation the local change needs, the second whether a larger update is needed for the quality target and fits the measured device budget, and the third whether the updates are valuable across devices. The flowchart's branch labels describe adaptation demand, not the Tier 1–3 device classes. Budget phones, held to the Tier 3 mechanism, sit in the lightweight branch, and the other Tier 1 and Tier 2 phones sit in the moderate branch. Wall-powered edge servers, which this phone fleet does not include, are candidates for the larger-update branch. The voice-assistant policy therefore spans

the lightweight and moderate branches, with budget phones held to the lightweight branch by a memory limit that the flowchart does not draw.

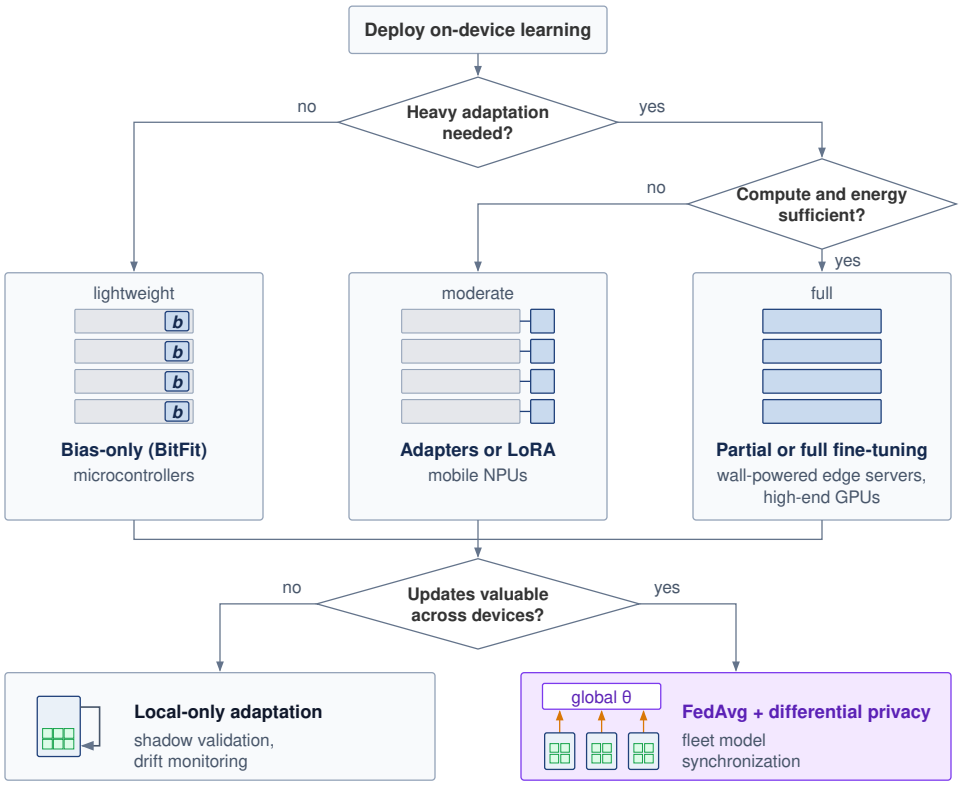

Figure 11.18: Tiered Adaptation Strategy: The flowchart routes a deployment by the adaptation the local change needs, then by whether a larger update is needed for the quality target and fits the measured device budget, then by whether updates are valuable across devices. Bias-only updates, adapters or LoRA, and partial or full fine-tuning are candidates for increasingly demanding changes on increasingly capable devices, subject to measured resource and quality budgets. The federated card names one common configuration; whether differential privacy is required depends on the deployment's threat model.

Checkpoint 11.4: Integrating the three pillars across device tiers

The chapter built three pillars for on-device learning: model adaptation (bias-only updates, adapters, full fine-tuning), data efficiency (few-shot, experience replay, compressed representations), and federated coordination (aggregation, client scheduling, secure aggregation). The decision framework in figure 11.18 maps each technique to device-class capability and coordination pathway. Test integration of the three pillars across the same deployment.

- ☐ Device-tier capability matching: For a fleet with microcontroller wake-word sensors, smartphones, and wall-powered edge servers, map each device class to the branch of figure 11.18 it can afford and state what fails when one technique is pushed across all three classes.
- ☐ Fairness in federated participation: When 90 percent of a federated round's compute comes from the top 10 percent of devices by capability, distinguish compute share from aggregation weight, identify the population-coverage risk, and select a scheduling or weighting control from the federated coordination pillar.

- ☐ Drift detection without labels: For a deployed personalization model with no ground-truth labels at the edge, describe how the device can detect enough local distribution shift to warrant a federated round using only signals available to the on-device inference path.
- ☐ Secure aggregation's participation floor: With a secure-aggregation floor of 100 devices and only 30 active devices in a given hour, determine the system response and identify which pillar must compensate.
- ☐ Observable signals and hidden state: Separate the signals a central operator can observe from the device state that must remain hidden, and locate the catastrophic forgetting risk in that split.

The hardest edge-learning failures appear when the pillars interact rather than when any one fails alone. An adapter can fit memory yet overfit sparse local data, a replay buffer can stabilize learning yet leak sensitive history, and a federated round can improve the global model while amplifying participation bias. These interactions decide when on-device learning is appropriate and when a simpler strategy is safer.

Centralized training can standardize hardware, curate data, and validate a model before release; a tiered policy instead ships to millions of devices whose runtimes, update histories, and connectivity differ, and it keeps changing the model after release. Shipping it turns each tier assignment into a release artifact that must be versioned, validated without central labels, and rolled back when an adapted model degrades.

11.10 Production Integration

A prototype can fine-tune an adapter on one Raspberry Pi under controlled conditions, where the engineer controls the input stream, inspects every loss curve, and reflashes weights by hand if an update diverges. Deploying that same adaptation policy across 50 million devices turns local learning into a production operations problem across devices whose runtimes, resource budgets, and update histories differ. The release contract must therefore govern not just a static set of initial weights, but what each device may learn in the field, how unlabeled local changes are gated, and how a trusted baseline is restored when adaptation fails.

11.10.1 MLOps integration challenges

In a traditional cloud pipeline, the release artifact is a single immutable model binary; once millions of devices begin updating adapters offline on different schedules, the initially released weights no longer describe the behavior running in the field. Integrating on-device learning into machine learning operations (MLOps), the practices that version, deploy, monitor, and roll back production models, expands the scope of release engineering because backbones, adapters, update policies, and local histories evolve independently across heterogeneous hardware. Existing integration tests, version registries, and monitoring remain necessary, but they must now track and bound the permitted adaptation process alongside the initial model.

An edge learning lifecycle splits five stages between the cloud and the device (figure 11.19). Centralized base-model pretraining (Stage 1), per-tier hardware compression (Stage 2), and staged over-the-air rollouts (Stage 3) originate in the cloud, while local adaptation executes on-device under the released update policy (Stage 4). Aggregate telemetry (Stage 5) supplies evidence for deciding whether cloud retraining is warranted. The telemetry needs its own privacy controls; aggregation alone does not establish a privacy guarantee, and a detected change must be distinguished from a quality regression before triggering a new training cycle.

11.10.1.1 Deployment pipeline transformations

A release that enables local learning includes an adaptation policy alongside model artifacts. For example, the voice-assistant policy of section 11.9 assigns bias-only updates to budget devices and different LoRA ranks to mid-tier and flagship phones. These are conditional assignments: model footprint, runtime support, and live admission tests determine whether a device may execute them.

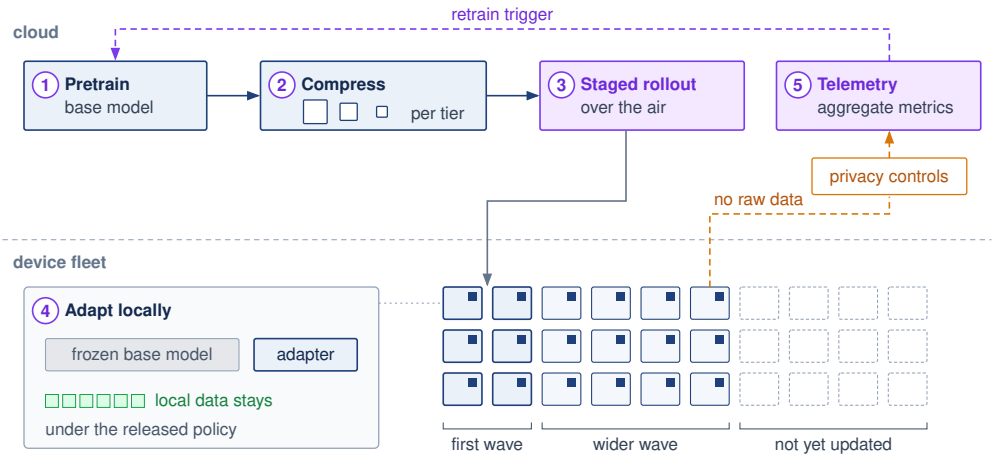

Figure 11.19: Edge MLOps Lifecycle Pipeline: The lifecycle links cloud pretraining and model compression with over-the-air deployment, on-device adaptation, and telemetry governed by explicit privacy controls. Keeping raw user data local does not by itself make the returned metrics private.

CI/CD must therefore validate the selection logic and permitted update behavior across device configurations in addition to the initial weights.

Versioning must preserve relationships among those artifacts. A local adapter belongs to a particular base model, training policy, and representation schema; a replay buffer may also depend on the encoder that produced its features. Federated round identifiers record which shared version produced a returned contribution. Updating the base therefore requires a decision to retain, migrate, retrain, or discard dependent local state. A newer version number alone does not establish compatibility, and rollback must restore a compatible set of artifacts rather than independently selecting older files.

11.10.1.2 Continuous training orchestration

Continuous training on devices combines fleet-level round or release schedules with local admission decisions. A coordinator can request work, but a device trains only when its resource and availability tests pass. The orchestrator must therefore track eligible work and compatible state without assuming simultaneous execution across clients. Because devices run different base versions at once and may reconnect after days or weeks offline, a returned update must identify its base version. The coordinator can then reject the update, apply a stated staleness policy, or request new training from the current base (section 11.7.5.1). Per-device adapters and replay buffers carry their own version tags and need not be merged into the shared model.

Resource-aware scheduling applies equation 11.2 at fleet scale, where federated rounds usually add idle, charging, and unmetered-network requirements on top of the local admission test. A daily energy cap bounds cumulative learning cost across many admitted windows. Admission also has a runtime half. Bias updates in early layers require backward computation through later operations; head-only adaptation confines that computation to the head. Either path competes with inference for compute, memory bandwidth, and thermal headroom on a shared SoC. Users judge the device by its quality of service (QoS), so the runtime gives inference priority and throttles or revokes a training window as foreground latency approaches its deadline. At the microcontroller extreme, measured update energy and available reserves determine how much learning a duty-cycled or harvested-power budget can support.[20]

When a representative central validation set is unavailable, federated evaluation can aggregate metrics computed on local held-out data. Update magnitudes provide an optimization diagnostic, but small changes can reflect stalled learning rather than acceptable quality. Participation audits must accompany outcome measurements so that apparent progress is not confined to an overrepresented cohort.

[20] **Microcontroller power budget**: In an illustrative operating point, 10 mW of active training for 10 seconds consumes 0.1 J. The number of updates that fit that session requires a measured time or energy per update. Battery lifetime additionally depends on usable capacity under the actual load, sleep power, radio activity, conversion losses, and duty cycle; nameplate capacity alone does not establish continuous-training lifetime.

11.10.1.3 Monitoring system evolution

Traditional monitoring aggregates metrics from centralized inference servers that see every request, whereas an adaptive fleet's monitor sees only what devices report about models that may differ on every device. Fleet telemetry for edge devices, including federated analytics that report aggregate or differentially private summaries, is one of the fleet release controls in section 11.10.1.5. Adaptation adds local model states that the operator may not observe centrally. Monitoring must distinguish a change in behavior from a change in quality. Unlabeled signals can identify divergence, while deciding whether the adapted model is better requires validated outcome evidence.

Drift detection presents additional challenges without access to ground truth labels. While traditional monitoring compares model predictions against labeled validation sets maintained on centralized infrastructure, on-device systems must detect drift using only the indirect local signals available during deployment. Confidence calibration compares output probabilities with eventual outcomes whenever labels exist; without verified outcomes, confidence shifts merely flag an altered score distribution. In parallel, input distribution monitoring applies statistical tests to detect when runtime feature distributions diverge from pretraining data without requiring labels. Task performance proxies use implicit feedback such as user corrections or session abandonment, though only after validating the causal link to task quality and accounting for how model suggestions alter subsequent user behavior. Finally, shadow baseline comparisons run a frozen base model alongside the adapted one; section 11.10.1.4 develops why that comparison measures divergence rather than quality.

Cohort evaluation must separate contribution from benefit. A device that never trains can still download an improved shared model, but its users' data may be poorly represented among the completed updates. Participation metrics reveal that coverage gap; per-cohort quality measurements establish whether it harms those users. A favorable global average cannot substitute for either measurement.

↪ Covariate shift, label shift, and concept drift, and why input monitors can stay silent under pure concept drift, are developed in section 14.4.1.

11.10.1.4 Validation strategy adaptation

Traditional validation measures model quality directly against a held-out test set with known ground truth, whereas on-device learning has to validate adapted models where their inputs live, within privacy and resource limits, and usually without labels. Shadow model evaluation maintains multiple model variants on each device and compares their behavior. Devices can run a frozen, last-known-good model beside a locally adapted version and measure prediction divergence, confidence shifts, latency, and resource use on the same input stream. Without labels or a validated outcome proxy, these comparisons detect change, not performance degradation: agreement can hide a shared error, and disagreement does not reveal which model is correct. A shadow comparison can therefore suspend an update for further evaluation, but automatic rollback requires labeled outcomes, a validated behavioral invariant, or a proxy whose relationship to task quality has been established for the deployment.

The comparator can therefore act only as a gate, never as a judge (figure 11.20). The live input stream is mirrored to the frozen baseline and the adapted candidate, and the comparator tracks their disagreement, confidence shift, and resource deltas. When disagreement exceeds a tolerance τ, the gate holds the candidate out of service; validated outcome evidence then either clears it or confirms the regression that justifies rollback.

Confidence-based gates provide another change signal when labels are unavailable. Even a calibrated model's confidence is interpretable only relative to the population and task on which calibration was measured. A drop can reflect distribution shift, ambiguity, or a changed score scale; stable confidence can coexist with overconfident errors. Systems may use confidence to trigger investigation or halt adaptation conservatively, but they should not call the adapted model worse without delayed labels, randomized evaluation, trusted user corrections, or another validated outcome signal.

Federated A/B testing randomly assigns devices to treatment and control adaptation policies, balanced across device tiers, and compares privacy-preserving aggregates of adaptation success rate, time to improvement, and final quality. A strategy that wins is then released through the staged rollout process that fleet operations owns. Because aggregated validation signals are themselves releases, they need data minimization, secure aggregation where appropriate, a declared privacy

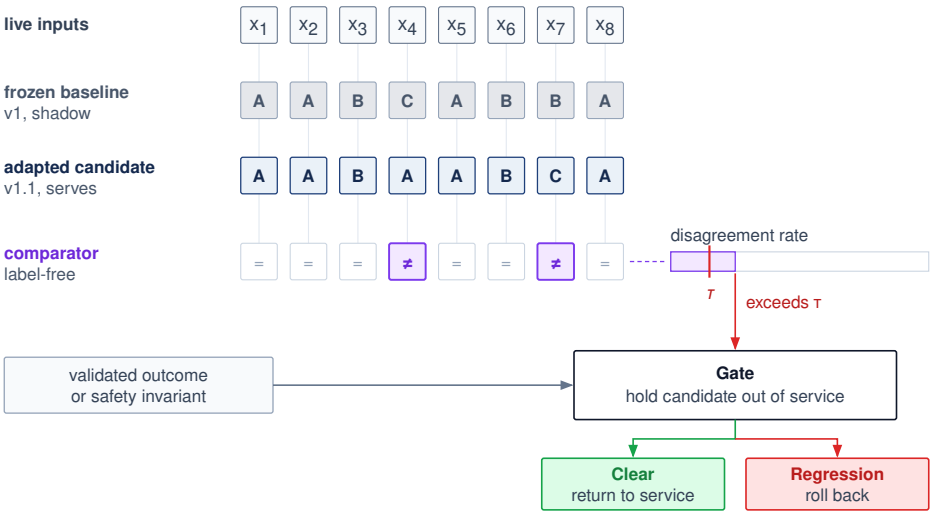

Figure 11.20: On-Device Shadow Validation: A frozen baseline runs beside the active model. Their disagreement, confidence, latency, and resource deltas can flag a change without labels; a rollback decision still requires a validated outcome signal or safety invariant.

unit, and privacy accounting across repeated rounds. Even then, a population summary can miss a failure confined to a few users. Local gates and population summaries therefore provide different evidence, and evaluating an adaptive fleet requires measuring adaptation benefit, resource cost, and cohort coverage together.

↪ Staged rollout, which widens a release's exposure only after each stage performs acceptably, is developed in section 12.5.5.

11.10.1.5 Fleet release controls

The preceding controls act on one adapted device or one federated round. Operating the whole population as one fleet adds three interacting constraints beyond those of any single device, so they form one operational design problem. Each constraint calls for its own platform control (table 11.9), including hardware-in-the-loop validation and federated analytics.

Table 11.9: Edge Fleet Controls: Edge operations couple rollout latency, device heterogeneity, and privacy-limited observability. The operations platform must reconstruct global fleet health from partial signals while preserving compatibility with many active model versions.

Edge fleet constraint	Why it appears	Platform control
Extreme version skew	Rollouts can take weeks or months because devices are offline, on low battery, or on restricted networks. At any time, 50 model versions may remain active.	Maintain backward-compatible data pipelines and serving contracts for models deployed months earlier.
Device-aware validation	Accuracy gates do not reveal whether a model exceeds a phone's memory budget or triggers thermal throttling on its system-on-chip (SoC); microcontroller-class targets below Archetype C tighten the budget to about 1 MB.	Add hardware-in-the-loop validation on physical or emulated target devices before promotion.
Privacy-limited telemetry	Raw predictions cannot stream back continuously because privacy rules and bandwidth costs constrain observability.	Use federated analytics: devices compute local statistics such as error rates or drift, then transmit permitted aggregate health signals under explicit privacy and retention controls.

Version skew and device-aware validation follow from the state an adapted fleet carries. A device that reconnects after weeks offline may still run a release several versions old, together with adapters

and replay buffers tied to that base, so data pipelines and serving contracts must stay compatible across every active version rather than only the newest. A cloud accuracy gate says nothing about the tests of equation 11.1 and equation 11.2, whose outcomes depend on each target's runtime, operator placement, and thermal behavior. Hardware-in-the-loop validation therefore runs the candidate release, including its update path, on physical or emulated devices from each tier, and promotion waits until those tests pass there.

Federated health signals inherit the participation bias that section 11.7.1 identified for federated training. A device reports its local error rate or drift statistic only when its admission policy allows the work and its connectivity allows the upload, so the devices that report most are the well-charged, well-connected, newer ones. The fleet's health dashboard then describes the admitted population rather than the deployed one, the same skew the participation selection principle (20) names for training. A staged edge rollout can look healthy because the devices most likely to fail, older hardware and thermally constrained phones, are also the least likely to report. The platform therefore tracks reporting coverage per device class and cohort alongside the health signals themselves, and holds promotion when a cohort's coverage falls too low for its signal to count as evidence.

11.10.1.6 Performance evaluation in dynamic environments

Static inference benchmarks characterize latency, throughput, energy efficiency, and accuracy for a fixed model, whereas on-device learning requires additional measurements of adaptation quality and training efficiency across successive updates to explain whether local learning is worth its resource cost. Sample efficiency measures the examples needed to reach a declared quality improvement under a fixed evaluation protocol. Gaining two percentage points from 100 local samples rather than 500 is an advantage when collecting useful samples is the binding constraint. Fewer samples need not mean faster or cheaper adaptation, because the more sample-efficient method may perform more work per example or retain more state; elapsed time, peak memory, and energy complete the comparison.

Because a sample-efficient update can still exceed hardware limits, device-envelope fit tests whether the improvement stays inside resource budgets. Memory-constrained convergence evaluates validation loss under a fixed RAM budget, such as "convergence within 256 KB training footprint," while energy-per-update measures millijoules consumed per gradient step. An illustrative background envelope of 500–1000 mW corresponds to 1.8–3.6 kJ/hour while active. Battery impact depends on duty cycle, usable battery capacity, foreground demand, and charging state. Time-to-adaptation extends the same idea to scheduling by measuring wall-clock time from new data to measurable improvement, including waiting for idleness, charging status, and thermal headroom rather than only raw accelerator throughput.

Even an update that fits the device envelope must still prove its value to the user, so personalization quality evaluates whether adaptation improves the intended behavior while retaining earlier capabilities. Per-user performance delta compares the adapted model with the global baseline on user-specific holdout data; acceptance requires a declared improvement threshold and reported uncertainty. Privacy must be assessed under the specified disclosure policy and guarantee, rather than treated as an unspecified unit of exposure. Catastrophic forgetting measures degradation on the original task after adaptation. Its acceptable magnitude is application-specific, and accuracy changes must distinguish percentage points from relative percent.

When devices coordinate through the federated protocols examined in section 11.6, evaluation expands from the device to the population. Communication efficiency relates quality improvement to total transmitted bytes; compression gains depend on the model, optimizer, data heterogeneity, and encoding (Konečný et al. 2016; Kairouz and McMahan 2021). Comparisons should also report bytes and elapsed time to the same quality target. Straggler policies change both waiting time and the cohort that contributes. Participation count, cohort coverage, and outcome uncertainty must therefore be evaluated together, without assuming that one participation threshold determines stability across all policies.

A system that infers quickly but adapts slowly, or adapts efficiently but ends at poor accuracy, still fails its deployment, so these training metrics and operational controls extend inference benchmarks and established MLOps practice to distributed, heterogeneous learning. Together they detect, hold, and reverse an adaptation that misbehaves, though they do not remove the underlying conditions

that make adaptive fleets misbehave, from heterogeneous hardware and non-IID data to slow drift and adversarial clients.

11.11 Engineering Challenges and Mitigations

A billion-user predictive text model that learns from live keyboard inputs can be poisoned if a coordinated group of malicious users repeatedly types a specific slur and those local updates enter the global federated aggregate. The release gates and runtime monitors of section 11.10 bound resource contention and flag gross regressions, and the tiered integration policy bounds where each pillar may act, yet neither erases the field conditions that let an adapting fleet degrade or conceal an attack. Those failure surfaces form a coupled chain across the deployment. Hardware and software fragmentation changes which devices can participate, skewed participation distorts the non-IID data distribution the fleet learns from, implicit feedback loops let local models drift silently between audits, and post-release weight changes strain regulatory traceability and safety fallbacks. Because each blind spot hides the next, engineering mitigations must progress from device and data heterogeneity through silent drift and compliance recovery to the combined control path that defends against coordinated poisoning (section 11.11.5).

11.11.1 Device and data heterogeneity management

An adaptation policy that passes lab tests on one reference smartphone can produce divergent weights, crash on unsupported kernels, or stall indefinitely across a production fleet whose OS versions, driver stacks, floating-point implementations, and charging habits differ. Heterogeneity determines which devices can train, which can validate, and which can only contribute lightweight signals. Tier assignment already gives each hardware class its own model format and update policy, so the management problem here is reproducibility across devices whose software stacks, connectivity, and power availability diverge.

Software heterogeneity is what makes even a correct assignment hard to reproduce. Devices may run different versions of operating systems, kernel-level drivers, and runtime libraries. Some environments support optimized ML runtimes like TensorFlow Lite Micro or ONNX Runtime Mobile, while others rely on custom inference stacks or restricted APIs. These discrepancies can lead to subtle inconsistencies in behavior, especially when models are compiled differently or when floating-point precision varies across platforms.

Connectivity and uptime add the scheduling dimension. Some devices are intermittently connected, plugged in only occasionally, or operating under strict bandwidth constraints. Others have continuous power and reliable networking, but still prioritize user-facing responsiveness over background learning. These differences complicate coordinated learning because update eligibility changes with the device's current state rather than with its hardware tier alone.

System fragmentation makes reproducibility and testing a fleet property rather than a lab property. With such a wide range of execution environments, consistent model behavior is difficult to guarantee and failures are difficult to reproduce. Monitoring, validation, and rollback therefore become more important, while also becoming harder to implement uniformly across the fleet.

The consequence is visible in a mobile-keyboard deployment. One device might provide 8 GB of RAM and a supported accelerator path, while another has 2 GB, a slower update path, and intermittent connectivity. These differences affect update duration and completion frequency. Adjusting schedules, model formats, and compression can broaden participation, but whether the improvement is equitable shows up only in outcome measurements across those device cohorts, because respecting each device's budget is not the same as delivering equal benefit.

11.11.2 Non-IID data distribution challenges

A federated round whose global validation loss improves steadily can still degrade next-word accuracy for an entire regional dialect or budget-handset cohort because no central shuffle can equalize non-IID client distributions. In production, non-IID data becomes an evaluation and control problem. Evaluation must cover the intended deployment population through a representative central test set or distributed held-out data, and aggregate quality needs cohort-level checks because a population average can conceal a persistent local failure. The mitigation should follow the observed

failure rather than the generic label "non-IID," and each of four common observations points to a bounded control and a release criterion (table 11.10).

Table 11.10: Non-IID Control Decisions: Non-IID data is not one scalar condition. The system should connect measured cohort skew or update conflict to a specific control and validate the result against the declared population objective.

Observed evidence	Likely failure	Control to test	Release criterion
Cohort metrics diverge while aggregate quality improves	Population underrepresentation	Stratified sampling or capped importance weights	Worst-cohort metric improves without unacceptable aggregate loss
More local steps increase update disagreement and rounds to target	Client drift	Reduce local steps, add proximal regularization (FedProx), incorporate control variates (SCAFFOLD), or adapt the server optimizer	Wall-clock time and bytes to target decrease on representative cohorts
One global model leaves stable, distinct clusters underserved	Multimodal client objectives	Clustered models or shared backbone with local heads	Each cluster has enough support and outperforms the global baseline
Heavy users dominate the aggregate	Objective-weight mismatch	Replace example weighting with user weighting or cap per-user contribution	Evaluation matches the declared user- or example-level objective

11.11.3 Failure modes and prevention mechanisms

Beyond skewing cohort averages, uncurated local learning can corrupt a model on a single device: a keyboard prediction system that trains on all user inputs without separating intended words from corrections can incorporate recurring typos as valid suggestions and progressively degrade its own predictions. Because such failures in on-device learning emerge slowly, locally, and often without centralized evidence (Kairouz and McMahan 2021), unbounded adaptation drift can pull a deployed model far from its intended behavior before aggregate telemetry moves. This risk becomes acute in health monitoring applications where gradual changes in user baselines could be learned as "normal," potentially causing the system to miss anomalies that a static model would have detected.

Beyond individual device drift, the participation bias described under client scheduling (section 11.7.1) can amplify itself over time. Users whose models improve less engage less, which further reduces their representation in later rounds (Wang et al. 2021). The failure is slow and invisible in aggregate metrics, which is why cohort-level evaluation must accompany every accepted round.

These systematic biases interact with data quality issues to create autocorrection feedback loops, particularly in text-based applications. When systems cannot distinguish between intended inputs and corrections, they may develop unexpected behaviors. Frequently corrected domain-specific terminology might be incorrectly learned as errors, leading to inappropriate suggestions in professional contexts. This problem compounds the drift issue, because models adapt to individual quirks and can simultaneously reinforce their own mistakes when users accept autocorrections without realizing the system learns from those interactions.

These failure modes connect individual drift, population bias, and data-quality degradation. Federated-learning surveys identify heterogeneity, participation bias, privacy, and optimization instability as deployment risks (T. Li et al. 2020; Kairouz and McMahan 2021). In response, successful systems can bound adaptation, stratify sampling, distinguish corrections from reliable labels, and use shadow comparisons to identify changes that warrant validated evaluation.

11.11.4 Compliance, security, and recovery risks in production

When a medical wearable, financial fraud detector, or vehicle controller adapts its weights after release, the deployed model diverges from the artifact that passed regulatory review. An unverified local update can invalidate the audit trail, absorb poisoned inputs, or violate data-erasure rules. In domains where compliance, auditability, and regulatory approval are required, on-device learning creates a direct tension between local autonomy and operational control.

A centrally managed training pipeline can record a common release's data, checkpoints, and evaluation results. On-device learning distributes the update history across local models and

observations that may remain private. Traceability therefore depends on what the adaptation protocol records: the base version, permitted update policy, validation result, and available recovery point. The operator can retain control over those boundaries without collecting every local training example.

This autonomy creates a validation gap. Evidence collected for the released model does not automatically establish the behavior of each adapted version. When inputs and update histories remain local, the deployment needs a defined set of local acceptance checks and a way to report failures within its privacy constraints. An adapted model can leave its validated operating range before aggregate monitoring detects the change, so the update policy must specify when to reject a candidate or restore the trusted baseline.

Delayed detection complicates recovery, especially while a device is offline. A trusted baseline must remain available independently of the adapted candidate, and the restore policy must account for compatible preprocessing, adapters, and runtime state. Restoring weights alone may leave the faulty update trigger or contaminated replay data in place. Recovery therefore includes deciding which adaptation state to retain, quarantine, or reset before learning resumes.

Every recovery path in this section assumes that something safe keeps operating while an adapted model is suspect or being restored. When perception and control loops rely on local models, that something is often a human supervisor, and a human fallback is not a safety mechanism unless the fallback itself is engineered as part of the system. A fatal 2016 crash under driver-assistance automation shows what happens when it is assumed instead.

War Story 11.1: The edge model with a human fallback (2016)

Context: A 2015 Tesla Model S was using Traffic-Aware Cruise Control and Autosteer when it collided with a crossing semitrailer near Williston, Florida, on May 7, 2016 (National Transportation Safety Board 2017). The driver remained responsible for supervision.

Mechanism: The NTSB attributed the crash to the truck driver's failure to yield and the car driver's inattention associated with overreliance on automation. The automation's operational design contributed by permitting prolonged disengagement and use inconsistent with the manufacturer's guidance.

Impact: The car passed beneath the semitrailer, and its driver died.

Fix: The NTSB issued safety recommendations addressing automation safeguards and oversight. For an edge-system designer, the corresponding requirement is to make operational limits and the supervision contract enforceable rather than assume a human will recover from every local failure.

Systems lesson: Fast local execution is insufficient if the combined sensing, automation, and human-supervision system lacks a reliable fallback. Although this vehicle did not learn on the device, the lesson carries over to edge learning. A fallback contract is validated against fixed model behavior, and local adaptation changes that behavior after release, so every accepted update must leave the operational limits and the validated recovery path intact.

Protecting the recovery path also requires securing the local state that triggers or feeds adaptation. Because on-device learning relies on device-specific data streams where central monitoring is limited, adversaries can tamper with replay buffers or inject poisoned examples during adaptation to push models into unsafe or biased states, while unauthorized access to stored feature embeddings or few-shot examples can leak personal information. The immediate design obligation is therefore to treat model state, replay buffers, and adaptation triggers as security-sensitive assets with local integrity checks, encrypted storage, and conservative update policies before adaptation begins.

Privacy and data-protection regulations compound these local state obligations. While keeping raw observations on the device reduces transmission exposure, adaptation still stores and processes personal sensor traces and behavioral logs locally, invoking requirements for user consent, retention limits, auditability, and the GDPR's right to erasure. Because a user's data can influence adapted weights and cached representations in ways that are difficult to trace and reverse individually,

compliant edge pipelines pair encrypted storage and retention windows with user-controlled reset mechanisms that purge local adaptation state and restore the validated baseline.

Distributed adaptation also complicates debugging and failure analysis (Kairouz and McMahan 2021). A deployment must assign responsibility for detecting regressions, suspending learning, restoring a baseline, and investigating failures before release. Local data drift, poor initialization, and faulty safeguards can produce similar symptoms, so the version and validation records must support diagnosis within the permitted disclosure limits. Autonomous execution does not remove the need for an accountable recovery process. Addressing these deployment and compliance risks requires tooling, protocols, and design practices that extend the ability to explain, bound, and reverse local adaptation (section 11.8.2) to external requirements for traceability, reproducibility, and user protection.

11.11.5 Poisoned updates and the combined control path

Poisoning, inversion, drift, and validation failures also occur in centralized ML systems, but local adaptation makes them harder to observe, attribute, and roll back across a heterogeneous fleet. The billion-user keyboard model poisoned by a coordinated group of users shows why. The malicious updates arrive as ordinary contributions from authenticated devices, secure aggregation hides each one from inspection, and the shift they cause can stay below what a population average reveals. Defending against such poisoned updates[21] is one of three obligations that local adaptation adds, alongside limiting what updates disclose and detecting local shift that population averages hide.

The edge architecture must therefore bind adaptation range, update integrity, privacy accounting, robust aggregation, validation evidence, and rollback into one control path. These mechanisms do not substitute for one another. Secure aggregation hides individual contributions, robust statistics limit some malicious influence, and validated outcomes establish whether behavior improved.

Traced through the slur attack, each control acts at a different point. Bounded local updates, checked before masking, cap how far any one device can push the shared model, although a coordinated group can still add many bounded pushes in the same direction. The masked sum keeps the coordinator from learning what each device typed, but it also keeps the coordinator from singling out the attackers' updates, so no per-client filter can remove them after masking. The change can surface in cohort-level evaluation and in shadow comparison against the frozen baseline, where the adapted model starts suggesting the slur, provided the evaluated inputs include the contexts the attack targets; a slur suggested only after a rare prefix can stay below the disagreement tolerance. When the comparison does flag it, the gate holds the candidate and fleet promotion waits; validated outcome evidence then confirms the regression and justifies rollback to the last compatible set of base model, adapters, and policy.

21 **Byzantine fault tolerance in federated learning (FL)**: Named after the Byzantine Generals Problem (Lamport et al. 1982), this property enables useful aggregation despite malicious or faulty participants; how many faulty participants a rule such as Krum or trimmed mean tolerates depends on that rule's own conditions, which section 7.3.3 states. These defenses can increase communication and compute overhead because they require extra validation, robust statistics, or redundant participation.

↪ Data poisoning as a threat model, including poisoned contributions to federated training, is developed in section 13.5.3.

11.12 Fallacies and Pitfalls

A team ports a cloud-trained recommendation model to a mobile app and assumes that successful inference also permits local fine-tuning. The first update can exceed the app's memory allowance as gradients and saved activations accumulate; a configuration that fits can still consume more energy than the background policy permits. These distinct failures motivate the following fallacies and pitfalls:

Fallacy: *On-device learning provides the same adaptation capabilities as cloud-based training.*

Teams expect local learning to match centralized training's model improvements, ignoring the device's resource constraints. Full-model training adds activations, gradients, and optimizer state to the weights; the configurations in section 11.3.1 require 4–12× the inference weights alone, depending on the model and execution schedule. Local datasets are typically small, biased, and nonrepresentative, while compute budgets remain orders of magnitude below cloud GPUs. A smartphone with 8 GB RAM and a sustained power budget of 2–3 W cannot replicate the adaptation achieved by data center systems with terabytes of memory and kilowatts of power. Effective on-device learning requires designing adaptation strategies that provide meaningful improvements within these constraints rather than attempting to replicate cloud-scale learning capabilities.

Pitfall: *Assuming federated learning automatically preserves privacy without additional safeguards.*

Practitioners believe that keeping data on local devices inherently provides privacy protection, ignoring what can be inferred from model updates. Gradient and parameter updates leak significant information about local training data through various inference attacks, while device participation patterns reveal sensitive information about users and activities. Stronger privacy preservation requires additional mechanisms. Differential privacy provides a tunable probabilistic bound on how much the output distribution can change because of one individual's data, while secure aggregation protocols prevent parameter inspection during server-side coordination through zero-sum cryptographic masking (figure 11.16) (Bonawitz et al. 2017). Data locality alone is insufficient for privacy protection.

Fallacy: *Resource-constrained adaptation always produces better personalized models than generic models.*

This belief assumes that any local adaptation is beneficial regardless of the quality or quantity of local data available. On-device learning with insufficient, noisy, or biased local data can actually degrade model performance compared to well-trained generic models (section 11.5). Small datasets may not provide enough signal for meaningful learning, while adaptation to local noise can harm generalization. Effective on-device learning systems must include mechanisms to detect when local adaptation is beneficial and fall back to generic models when local data is inadequate for reliable learning.

Pitfall: *Ignoring heterogeneity across different device types and capabilities.*

Teams design on-device learning systems assuming uniform hardware capabilities across deployment devices. Real-world deployments span diverse hardware with varying computational power, memory capacity, energy constraints, and networking capabilities. Available memory alone spans roughly 65,536× between 256 KB microcontroller SRAM and 16 GB flagship phones (section 11.7.5.2), and supported operators and power envelopes differ as sharply. Clock frequency or peak ratings from unrelated architectures do not convert into a common training-throughput ratio, so the comparison needs measured update cost under matched tasks and declared operating states. A learning algorithm that works well on high-end smartphones may fail catastrophically on resource-constrained IoT devices. Federated learning algorithms must account for heterogeneous clients and availability (Kairouz and McMahan 2021; Bonawitz et al. 2019), using selective participation, compression, and tiered aggregation strategies when devices have substantially different measured update costs.

Pitfall: *Underestimating the complexity of coordinating learning across distributed edge systems.*

Many teams focus on individual device optimization and treat coordination as a stateless averaging server, as if local updates arrived regularly from comparable devices. At fleet scale the aggregator is one component of a control loop that must know which devices are eligible, which model version each holds, and whether a returned update is safe to promote. A production system selects about 130 percent of its target cohort (Bonawitz et al. 2019) because a substantial fraction of selected devices never report. Edge orchestration (section 11.7.5) must also handle intermittent connectivity, varying power states, time zones, and version skew across devices that return after days offline. A design that budgets only for the averaging step leaves round deadlines, partial participation, and state reconciliation unowned, although those parts of the loop decide whether each round completes with a representative cohort.

11.13 Summary

Edge intelligence extends the fleet to devices whose memory, power, and availability the operator cannot standardize. Sensors, microcontrollers, and smartphones contribute under different resource budgets, so a locally feasible learning policy also determines which parts of the population can participate. Archetype C (Federated MobileNet) makes this concrete. Local adaptation spends the compute axis of C^3 inside each phone's power and memory envelope. Adapters and update compression shrink the wide-area transfer in the Fleet Law's communication term. Over-selection with round deadlines caps the Fleet Law's synchronization term at the price of which devices are represented.

Three levers make learning possible under those constraints, and each acts on a different part of the same budget. Model adaptation shrinks trainable state and, when placed well, the saved

activations that dominate peak memory. Data efficiency raises the signal extracted from each scarce example and protects earlier behavior with bounded replay. Federated coordination lets compatible updates from many devices improve one shared model without centralizing raw data. None of the three moves the device's thermal, energy, or bandwidth limits; each decides how much useful learning fits inside them.

At runtime, those realities become an admission decision rather than a general preference to run "when idle." Background adaptation may execute only when its predicted energy and sustained power fit the current device slack and profiled worst-case interference preserves the foreground deadline. Because charging status does not eliminate heat or contention, every admitted window remains revocable as the device state changes.

Local execution can meet latency, data-locality, and offline requirements, but those requirements alone do not justify local training. Adaptation must improve a validated outcome within its resource budget; federation must then preserve update compatibility, control disclosure, and evaluate the population actually represented by completed updates.

Key Takeaways: Physics sets the budget

- **The edge reverses the fleet**: Data centers fight to keep many accelerators busy; edge devices fight to learn under batteries, kilobytes, thermal throttling (the device-scale counterpart of principle 3), and intermittent links. The same Compute, Communication, Coordination constraints must account for device eligibility as well as aggregate throughput.
- **Bandwidth beats advertised TOPS**: One update step of the chapter's 10M-parameter example moves about 800 MB, 20× the weight bytes read by an inference, across a phone memory interface roughly 44–65× narrower than data-center HBM. That is the same ceiling that bounds memory-bound decode (principle 18). Quantization and locality help, but measured update speed depends on bytes moved and sustained bandwidth.
- **Learning multiplies the footprint**: Full-model backpropagation can require 4–12× the memory occupied by inference weights alone in the chapter's configurations, depending on optimizer state, activations, and precision. Reducing trainable state helps only when the frozen model and remaining execution state also fit.
- **Background learning is conditional work**: Admit an adaptation window only when equation 11.2 preserves energy, thermal, and foreground-latency budgets under worst-case measured interference. Charging alone is not sufficient.
- **Admission selects the population**: Passing the admission test correlates with hardware, charging habits, and connectivity, so the shared model learns from the admitted devices rather than the deployed fleet (principle 20). Per-cohort participation and quality must be measured and corrected; loosening admission spends the budgets it exists to protect.
- **Each technique spends a different budget**: Bias-only updates, LoRA, sparse layers, few-shot data reuse, and experience replay each spend different memory, energy, and privacy budgets. Federated coordination adds population learning only when cumulative uploads stay below the data they replace, enough clients survive to clear secure aggregation's participation floor so individual updates stay hidden, and shared state stays version-compatible.
- **Heterogeneity is normal operation**: Edge fleets span microcontrollers, phones, gateways, and NPUs with uneven availability, performance, and non-IID data. Stragglers, client drift, participation bias, and privacy-preserving observability are protocol design constraints, not deployment cleanup. Once models learn after release, the release artifact is the adaptation policy, versioned against a base model that rollback can restore.

An edge learning policy determines both the work a device may perform and the population the fleet learns from. Tightening a resource budget can protect a user's foreground experience while excluding that user's data from shared improvement. The useful design target is therefore validated

adaptation across the intended population, with local feasibility and fleet representation evaluated together.

What's Next: From devices to operations

Once a model continues learning after release, what the fleet ships is no longer a fixed set of weights but an adaptation rule versioned against a base model, admitted device by device, validated through distributed outcome evidence when central labels are unavailable, and judged by which cohorts its completed updates represent. Production organizations run hundreds of models at once, many sharing features, data pipelines, and release machinery, so a change made for one can silently break another. Keeping that portfolio versioned, observable, and recoverable is a coordination problem in its own right.

↪ The control plane that versions, releases, monitors, and rolls back a portfolio of models is developed in chapter 12.

12

ML Operations at Scale

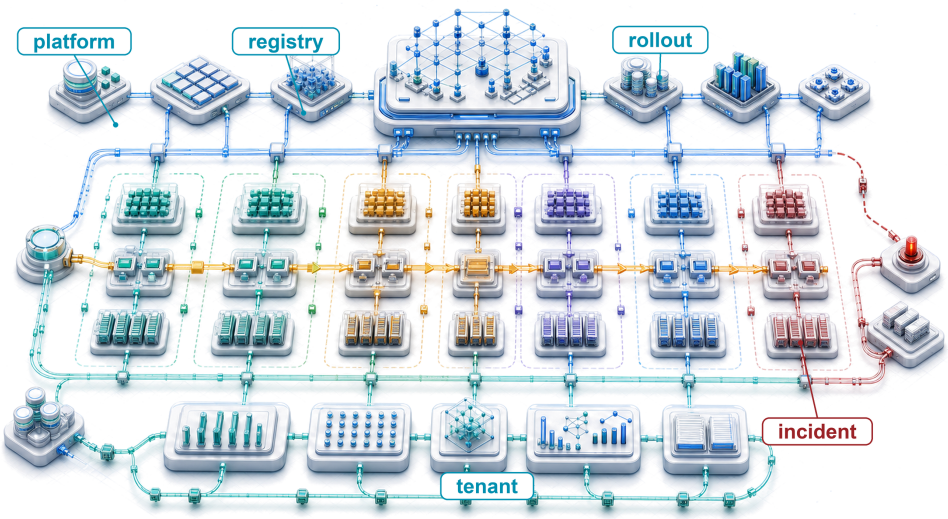

12.1 Why a Model Fleet Needs a Control Plane
12.2 The N-Models Problem
12.3 Governance of State and Change
12.4 Feature Store Operations
12.5 CI/CD for ML at Scale
12.6 Monitoring at Scale
12.7 Production Debugging and Incident Response
12.8 Platform Engineering
12.9 Fleet Economics and Utilization
12.10 Organizational Operation and Learning
12.11 Case Studies
12.12 Fallacies and Pitfalls
12.13 Summary

Purpose

Why do practices that work for managing one model collapse when organizations deploy hundreds?

One model is a project. A hundred models form a system of systems whose dependencies and failures cannot be managed one model at a time. A data-pipeline change may affect models built by several teams, while a deployment failure may require coordinated rollbacks across interconnected services. The team making a change may not know every consumer, and the team seeing a regression may not control its cause. Practices that work when one engineer knows the entire system become unreliable when ownership, data, and release schedules are distributed. Adding faster machines does not resolve that uncertainty. A fleet can have spare compute capacity and still be unable to ship safely because nobody can establish which models depend on a changed input or who may authorize recovery. Shared infrastructure can reduce this repeated coordination, but it also concentrates dependencies and introduces its own operating cost. Centralizing every decision can turn the platform team into a bottleneck; leaving every team to improvise makes compatibility and accountability difficult to sustain. The engineering task is to make routine changes repeatable while preserving evidence about their effects and clear authority over their consequences. Growth becomes manageable when adding a model does not require rebuilding the surrounding operational process, yet its owners remain responsible for acceptable behavior. In C^3 terms, fleet operations makes coordination a designed system constraint alongside the compute and communication resources that models consume.

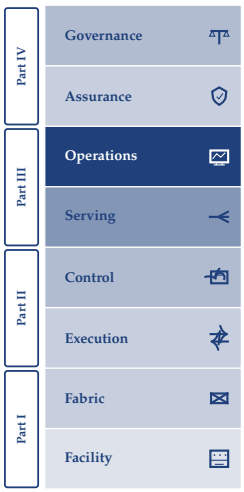

Learning Objectives

- Explain why operational load grows with coupling through shared features, pipelines, and capacity rather than with model count
- Design registries, release gates, and evaluation contracts that preserve lineage, dependency safety, and rollback confidence
- Design feature-store operations that manage freshness, ownership, compatibility, and training-serving skew
- Size staged rollouts and production experiments whose sample size, stopping rule, and interference assumptions bound user exposure
- Diagnose production incidents using monitoring hierarchies and multidimensional reliability budgets across the fleet
- Calculate platform return, utilization break-even, and lifecycle total cost, then attribute fleet spend to accountable owners
- Evaluate platform organization and lifecycle choices using adoption, versioning, and retirement evidence

12.1 Why a Model Fleet Needs a Control Plane

Consider a team of five engineers maintaining a single recommendation model. When the model's accuracy drifts as live data changes, they manually retrain it. When the API latency spikes, they manually scale the instances. Now, scale that same team to support five hundred models across dozens of product surfaces. The same engineering team cannot sustain that growth by performing machine learning operations (MLOps), the practices that version, deploy, monitor, and roll back production models, by hand for every model. Shared controls must preserve human accountability while automating routine coordination. The fleet becomes the unit of capacity planning and investment. Versioned model and feature state makes changes reviewable, release gates bound their exposure, and monitoring and incident review feed production evidence back into those controls.

In the fleet stack of figure 1.4, the Operations row belongs to the Serving band, just below the Governance band, because it operates models that are already deployed; the lower Reliability and Control row schedules and recovers individual jobs. The Operations row acts as the fleet's control plane, the dashboard, steering, and maintenance system that keeps physical infrastructure, training systems, serving paths, edge deployments, and governance rules from drifting apart. Once models span data centers, serving platforms, and heterogeneous edge fleets of adaptive devices, reliability depends less on any single model and more on the operational machinery that keeps the whole fleet observable, coordinated, and recoverable. Adaptive edge fleets sharpen the requirement, because there the operator ships a versioned adaptation policy rather than fixed weights, and that policy stays safe only while telemetry keeps it observable and rollback keeps it recoverable. Those controls must operate as shared infrastructure for every model the fleet runs.

↩ Adaptive edge device populations and the constraints that shape them are established in chapter 11.

Shared dependencies make local operational choices fleet-wide concerns. An upstream data change can affect several models, concurrent deployments can contend for the same capacity, and independent alerts can obscure a common failure. Training and serving also compete for expensive infrastructure whose idle capacity still incurs cost. The platform must coordinate those interactions while preserving each team's responsibility for model behavior.

The organizational response is platform thinking. Rather than treating each model as an independent system with its own infrastructure, platforms provide shared services that amortize operational costs across the entire model portfolio. Feature stores[1] share feature definitions and values across models. Their contract is that training and serving use compatible feature meanings, with historical training values restricted to information available at the prediction time. Unified deployment pipelines ensure consistent rollout practices. Centralized monitoring aggregates signals across models to detect system-wide issues and enable capacity planning. The challenge is designing, implementing, and operating these platforms so they scale with the portfolio rather than against it. This is the C^3 coordination axis at its widest. The Fleet Law charges synchronization within one

[1] **Feature store**: A centralized repository that manages the computation, storage, and serving of ML features. The core systems problem it solves is training-serving skew; section 12.4 develops that failure mode and the platform invariant that contains it.

distributed step, while fleet operations pays the coordination cost of every release, rollback, and shared dependency across the whole portfolio.

12.2 The N-Models Problem

The five engineers of the opening example do not acquire five hundred models overnight, and the path from one model to many explains where a control plane becomes necessary. An organization's first homepage recommender, search ranker, and fraud detector can each run on a bespoke, team-local pipeline for data processing, training, validation, and deployment that lets each group optimize for its immediate requirements. As the fleet grows, those independent pipelines begin sharing features, embeddings, and GPU pools, until one uncoordinated upstream change breaks five downstream models at once.

As the number of models grows, shared dependencies create **operational complexity** faster than model count alone suggests. Dense interactions produce superlinear coordination work unless strong ownership and interface boundaries contain them. Between one model and a hundred, deployment coordination moves from a team-local task to the critical path, and shared data dependencies grow from a single pipeline into a dense graph. Monitoring, on-call scope, utilization, debugging, and governance change with them (table 12.1).

Table 12.1: Operational Complexity Growth at Scale: Illustrative challenges across 1, 10, and 100 models. Even one model requires deployment coordination and depends on its data pipeline. Shared dependencies can make releases a fleet-wide critical path, dashboards require aggregation, and debugging span team boundaries; actual complexity depends on coupling and ownership.

Operational aspect	Single model	10 models	100 models
Deployment coordination	Team-local	Ad hoc	Critical path
Shared data dependencies	Within pipeline	Some overlap	Dense graph
Monitoring dashboards	1	10	Unmanageable
On-call rotation scope	Single team	Multiple teams	Organization-wide
Infrastructure utilization	Often idle	Moderate sharing	Efficiency critical
Debugging complexity	Local	Cross-team	Distributed tracing required
Governance	Team policy	Diverging policy	Organization-wide automation

Napkin Math 12.1: Shared versus dedicated accelerator pools

Problem: A platform team manages a fleet of 100 GPUs. Under dedicated per-team quotas, average idle time is 70 percent. Moving to a multi-tenant ML platform that shares resources across teams and uses idle training GPUs for inference reduces aggregate idle time to 30 percent. What hardware cost does the platform save?

Math: For a fixed active workload, required hardware is inverse to utilization, while useful work per GPU is proportional to utilization. The problem states idle time, so the calculation first converts it to utilization.

1. **Utilization**: Utilization is one minus the idle fraction, so dedicated quotas run at 0.30 and the shared platform at 0.70.
2. **Efficiency Gain**: 0.70 (shared utilization) / 0.30 (dedicated utilization) = 2.33× more work per GPU.
3. **Hardware Reduction**: $1 - (0.30 / 0.70) = 57.1$ percent.
4. **Annual Savings**: 57 percent of \$1.75M budget ≈ \$1,000,000/year.

Systems insight: Multi-tenancy acts as an infrastructure multiplier. Under the assumed utilization improvement, breaking down resource silos reduces required hardware by 57 percent for the same workload; with the same hardware budget, it raises useful work from 30 to 70 active GPU-equivalents. In the machine learning fleet, *statistical multiplexing* pools

demands whose peaks do not fully coincide, allowing teams to share capacity that separate allocations would leave idle. Realizing these savings requires enough headroom for coincident peaks and failures, together with isolation that preserves each workload's service contract.

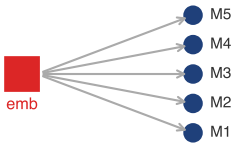

One upstream embedding update can degrade every dependent model.

Per-model validation does not establish compatibility across a dependency chain. When Model A depends on features computed by Pipeline B, which uses embeddings from Model C, a change to one component can affect its downstream consumers. A seemingly innocuous update to Model C's embedding layer might shift the feature distributions that Model A depends upon, degrading its performance even though Model A itself has not changed. This cascading interdependence makes the interfaces between models part of the operational contract.

Systems Perspective 12.1: The complexity explosion

As a fleet accumulates shared dependencies, system-level coordination can become the binding problem. A new model adds its own operating cost, but it can also change feature pipelines, release schedules, and contention for infrastructure used by existing models. The marginal cost therefore depends on how the new model couples to the portfolio rather than on how many models the fleet already contains.

A sketch of operational load across production model counts, drawn on logarithmic axes with one assumed growth shape per source of work, shows how coupling outruns a fixed team (figure 12.1). Alerts grow linearly because each model brings its own fixed set of monitored metrics. The sketch assigns deployment coordination an intermediate $N_{\text{models}} \log N_{\text{models}}$ growth shape; this is an illustrative assumption rather than a release-scheduling complexity bound. A dense dependency graph motivates the quadratic conflict curve because N_{models} models form roughly $N_{\text{models}}^2/2$ possible pairs; the recorded dependencies determine which pairs can actually conflict. Under these shapes the total crosses the team-capacity ceiling near 50 models. The shapes are a sketch, not measurements, so the crossing marks where repeated work outgrows a team rather than a threshold a real organization can read off the plot. Measured dependency density, engineering effort, and remediation cost decide when platform engineering pays.

Once operational load exceeds team capacity, continued portfolio growth requires reducing repeated work, simplifying dependencies, or expanding the team's capacity. Shared platform engineering addresses the repeated work of deploying, monitoring, and debugging models. Its economic case begins with identifying which of those responsibilities can be shared and which must remain with model owners.

12.2.1 Single-model vs. platform operations at scale

Separating shared platform duties from model-owner responsibilities starts with seeing where single-model habits fail in a coupled portfolio. A single-model team can promote a release through a canary (a small slice of live traffic served by the candidate) and roll back in isolation. In a shared fleet, pushing an updated upstream model can silently shift inputs for downstream rankers, exhaust shared GPU headroom during concurrent rollouts, or leave dependent services broken if one model rolls back alone. Platform-scale deployment must preserve compatibility with dependent models at every intermediate state. Compatible interfaces can permit independent updates; a changed feature representation may instead require a coordinated migration. Rollback follows the same dependency contract, because reverting one model may require restoring compatible features or dependent models. Concurrent deployments also compete for GPU memory and network bandwidth, making resource headroom part of the release decision.

For recommendation systems, this complexity appears along the request path. Candidate retrieval and ranking models execute alongside diversity filters and business rule layers, sometimes in sequence and sometimes in parallel. Updating a component requires identifying which consumers depend on its outputs and testing whether the complete path still meets its quality and latency targets.

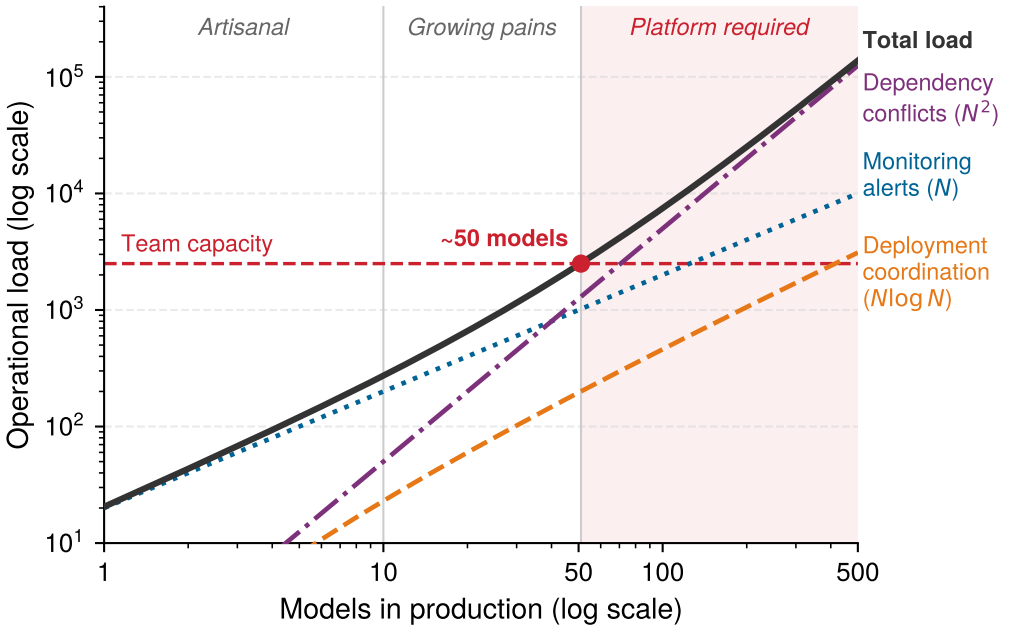

Figure 12.1: The N-Models Complexity Explosion: The curves show monitoring alerts growing linearly with model count, deployment coordination growing as $\mathcal{O}(N_{\text{models}} \log N_{\text{models}})$, and dependency conflicts growing quadratically as models share features and data sources. Under these assumed shapes, the total operational load crosses team capacity around 50 models, sketching the transition from artisanal model management to platform-required operations.

Monitoring requirements evolve similarly. Model-specific metrics such as prediction accuracy, inference latency, and data drift remain necessary, but separate dashboards do not expose shared causes by themselves. When many models depend on the same feature pipeline, a single upstream failure can appear as many unrelated model alerts. Responders need a portfolio view that connects those symptoms before investigating individual models.

Platform monitoring must therefore aggregate across models while maintaining the ability to drill down into specifics. This requires hierarchical metrics that run from business outcomes through portfolio and model signals down to infrastructure health, the structure section 12.6.3 develops. The same contrast between one model and a portfolio also changes what each release must prove, because workloads differ in how fast they change and how much harm a bad version can do.

12.2.2 Release policies across model archetypes

Applying a single release policy across a fleet either starves a fraud detector that needs hourly updates and second-scale rollback as attackers adapt, or rushes a 70B–405B generalist large language model (LLM) into production before multi-domain evaluation catches a capability regression. Different deployment contexts impose different operational contracts, and the archetype taxonomy in section 1.6.1 helps interpret the illustrative policies in table 12.2. Release frequency follows data change and validation cost; rollback targets follow the harm and reversibility of a bad release. Model family alone determines neither: LLMs may need long observation windows while still requiring rapid traffic rollback, and fraud detection needs fast containment even when a new model takes longer to validate.

Table 12.2: Illustrative Release Policies: Example cadences and rollback targets for several deployment contexts. These are planning assumptions, not requirements implied by architecture. Rollback means routing to a validated version; the observation period needed to approve a new version can be much longer.

Model type	Example update cadence	Deployment pattern	Primary risk	Rollback speed
Archetype A (GPT-4/Llama-3)	Monthly to quarterly	Staged, careful	Quality regression, safety	Minutes with standby capacity
Archetype B (DLRM at Scale)	Daily to weekly	Shadow, interleaving	Engagement drop	Minutes
Fraud Detection	Hourly to daily	Rapid with instant rollback	False negatives	Seconds
Vision (Classification)	Weekly to monthly	Canary	Accuracy regression	Minutes
Search Ranking	Daily	A/B with holdout	Relevance degradation	Minutes

The table mixes two lighthouse archetypes with ordinary model families because release policy follows the workload, not the label. Archetype C (Federated MobileNet) is absent because its edge fleet follows release controls of its own (section 11.10.1.5). The deployment patterns name release controls developed later in the chapter. Shadow deployment runs the candidate on copied traffic without returning its outputs to users, and interleaving merges two rankers' results into one list so each user compares them directly (section 12.5.3).

The table separates release cadence from containment speed. Large-model evaluation can require extended shadow deployment, human review, and staged exposure because a quality regression may be subtle. Keeping the previous model warm can nevertheless permit fast traffic rollback. The cost of that standby capacity belongs in the release policy rather than being hidden behind an hours-long rollback assumption, and the cost of regression becomes concrete when one model must preserve every capability at once.

Lighthouse 12.1: Archetype A (GPT-4/Llama-3): Cost of regression

Context: Archetype A (GPT-4/Llama-3) represents ultra-large generalist models (70B–405B parameters) serving millions of concurrent requests across diverse tasks. Released models span that whole range; where this chapter needs one concrete instance, as in the fleet cost examples of section 12.9, it uses the book's reference instance, the GPT-3-class dense model of section 1.6.1.1.

Mechanism: The "Generalist's Dilemma" occurs when targeted fine-tuning (e.g., optimizing Python code generation) causes catastrophic forgetting, the loss of previously learned capabilities, in unmonitored domains (e.g., creative writing or reasoning). Evaluating candidate weights across 10,000+ benchmark prompts adds the iron law's compute time, $O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$ for the evaluation's total operations O, to every release. Prompt count alone does not determine that time or its GPU-hours; input and output lengths and model size set O, and achieved throughput sets η_{hw}.

Impact: A silent degradation in secondary metrics forces a costly emergency rollback, and the warm standby held for that rollback competes with the serving capacity needed to hold the per-request latency target $L_{\text{lat,SLO}}$.

Systems lesson: Release gates for large generalist LLMs must combine policy-based safety checks with broad capability benchmarks such as Massive Multitask Language Understanding (MMLU), a multiple-choice knowledge test across many subjects, and HumanEval, a code-generation test. Their evaluation time and accelerator-hours belong in the deployment schedule and budget rather than being treated as a zero-cost pre-flight check.

Recommendation systems face continuously changing user preferences and content inventories (Steck et al. 2021; Gomez-Uribe and Hunt 2015), so their operational urgency includes feature freshness, the delay between a user action and its reflection in the model's inputs, which section 12.4.2

measures and budgets. Recommendation ensembles must also preserve compatible behavior as their components update. Fraud detection adds an adaptive adversary that probes the model and shifts tactics once detected, increasing the value of prompt containment and monitoring for sudden feature shifts. Shadow scoring and outcome evidence still have to establish whether an update reduces fraud without unacceptable false positives; urgency alone does not validate a new model. LLM releases face a broad set of capability and safety risks that require evaluation before exposure (Bender et al. 2021). These differences determine the evidence, update cadence, and recovery policy each workload needs. A shared platform is valuable when it enforces those distinct contracts without making every team rebuild the operational machinery.

12.2.3 Quantifying platform economics

Enforcing those workload-specific release policies without forcing every team to rebuild the machinery requires a shared platform, which raises an immediate budget question. Standing up a dedicated ML platform team for a five-model organization can cost twice as much as the manual toil it replaces. Delaying that investment past fifty models, by contrast, leaves product teams buried in repetitive deployment and pipeline maintenance. The economic case for platform operations rests on comparing the cost of fragmented approaches against the returns from shared infrastructure. Equation 12.1 defines the platform benefit-cost ratio (BCR) as the ratio of engineering time savings across all models to total platform cost:

$$\text{BCR}_{\text{platform}} = \frac{N_{\text{models}} \times T_{\text{saved}} \times C_{\text{engineer}}}{C_{\text{platform}}} \tag{12.1}$$

where N_{models} represents the number of models benefiting from the platform, T_{saved} is the engineering time saved per model per period, C_{engineer} is the fully-loaded cost per engineer hour, and C_{platform} is the total platform cost including development, infrastructure, and maintenance.

A ratio of one is break-even; net return on investment is this ratio minus one. Engineering hours valued at their loaded cost measure recovered capacity, not automatic payroll savings. Costs and benefits must cover the same period and include ongoing platform operation.

The equation makes the investment boundary explicit. For an organization with five models, platform cost may exceed the value of recovered effort; unusually costly recurring work could reverse that comparison even at a small model count. Holding per-model savings constant makes the numerator grow linearly with N_{models}. Whether the ratio improves depends on how development, support, and infrastructure costs grow as additional models adopt the platform.

> Napkin Math 12.2: Platform benefit-cost ratio
>
> **Problem**: An organization manages 50 models. A centralized ML Platform team costs $120,000/month. If the platform saves each model team 20 hours of manual toil per month, is the platform investment profitable?
>
> **Math**:
>
> 1. **Gross Monthly Capacity Value**: 50 models $\times$ 20 hours/model $\times$ $150/hr = $150,000.
> 2. **Net Monthly Benefit**: $150,000 (Savings) - $120,000/month (Cost) = $30,000.
> 3. **Benefit-Cost Ratio**: $150,000 / $120,000/month = 1.25.
>
> **Systems insight**: Platforms cross a scaling threshold under a declared cost model. At 50 models, this platform earns a 25 percent return on its modeled cost. However, if the organization only had 20 models, the capacity value would be only $60,000, a 50 percent loss relative to the platform team's salary. This comparison holds the team cost and per-model savings fixed. An investment decision must also include migration, infrastructure, and ongoing support costs, then establish how the recovered engineering time will create value.

12.2.3.1 From artisanal to industrial operations

The notebook prices one portfolio at one model count. The investment decision also needs to know where break-even falls and which inputs move it.

Plotting the benefit-cost ratio against production model count for two annual platform budgets locates the break-even crossing, where each linear benefit ray climbs out of the shaded investment phase into the net-positive returns phase (figure 12.2). With a fixed $100K benefit per model per year, a $2M/year platform achieves parity at the 20-model marker, whereas a $5M/year enterprise platform requires 50 models to reach the break-even line. At 100 models, the $2M platform ascends to a benefit-cost ratio of 5, equivalent to a net ROI of 400 percent. The linear curves assume constant marginal benefit and fixed platform costs; dependency growth, diminishing savings, or higher operating costs change the result. The $5M platform's 50-model crossing is also a different quantity from the 50-model crossing in figure 12.1: that one marks where repeated work outgrows a team, whereas this one marks where a budget pays back. The figure's budgets are annual running costs, while a one-time build amortized over three years charges only a third of its cost to each year. In either form, the break-even count moves with the two inputs a platform team controls, shifting left as savings per model rise and right as platform cost rises.

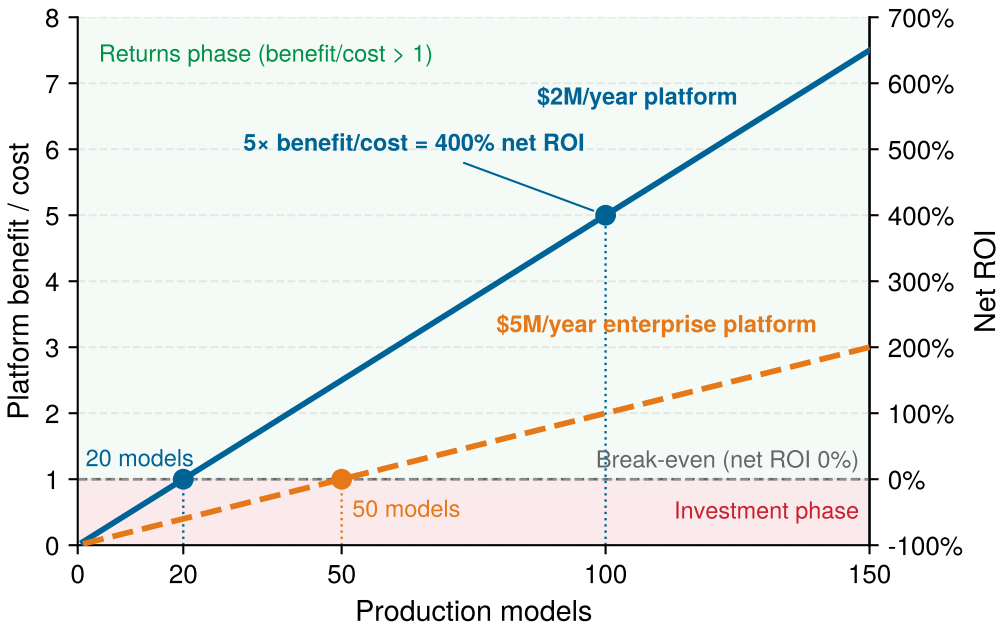

Figure 12.2: The Platform Benefit-Cost Threshold: Assumed annual benefit divided by annual cost for two platform budgets. Break-even occurs at a ratio of one. With fixed per-model benefit, the ratio grows linearly; a ratio of five corresponds to a net ROI of 400 percent, not 500 percent.

Checkpoint 12.1: Platform ROI break-even

Equation 12.1 expresses the platform benefit-cost ratio as $N_{\text{models}} \times T_{\text{saved}} \times C_{\text{engineer}}/C_{\text{platform}}$. In figure 12.2, the two benefit-cost rays show this linear growth under fixed savings assumptions and where each platform budget crosses break-even. Apply the formula to three concrete decisions.

- ☐ **Platform ROI**: for an organization with 30 models, each saving 200 engineer-hours per year at $150/hour, calculate the benefit-cost ratio and net ROI of a $2M/year platform and a $5M/year enterprise platform.
- ☐ **Break-even count**: using the same per-model savings ($30K/year, which is 200 hours at $150/hour), determine the model count where each platform breaks even, compare it to

the 20- and 50-model parity markers plotted in figure 12.2, and explain how the figure's stated $100K per-model benefit moves those thresholds.

- ☐ **Build now vs. wait**: for an organization with 8 models debating whether to build a $2M/year platform, use the linear ROI structure of equation 12.1 past break-even to state the strongest argument for building now and the future-growth assumption required for that argument.

The benefit-cost ratio takes the per-model savings T_{saved} as an input. In practice those hours come from specific recurring work, such as unversioned data, brittle pipelines, and manual deployment toil, that a shared platform can retire. Measuring that work is the job of technical-debt accounting. The hardware side of the cost accounting, where utilization decides whether shared capacity earns its cost, returns in section 12.9 after the operational machinery that consumes that capacity is in place.

12.2.4 Quantifying and managing ML technical debt

Those recurring hours accumulate whenever unvalidated 800-line configuration files, duplicated preprocessing scripts, and ad-hoc monitoring setups spread across dozens of models, stretching a two-day training run into a six-week release. Technical debt at fleet scale becomes a prioritization problem when recurring maintenance competes with new development. The platform must identify dependencies that slow deployments, raise incident rates, or consume engineering time, then compare the expected benefit and cost of repairing them.

ML debt falls into four categories according to where the recurring work originates. Data debt begins upstream in inputs that are unversioned or weakly validated. Configuration debt accumulates in duplicated, unvalidated settings that every release must review by hand. Closer to the artifact, model debt arises from the coupling around a model, including undeclared consumers and **glue code**, the supporting code written to move data into and out of general-purpose libraries and between pipeline stages. Glue code binds a model to one pipeline's formats. Infrastructure debt surfaces in brittle pipelines, manual deployment, and drifting environments. The categories locate where the work originates; common effort and risk measures make candidate repairs comparable across the fleet.

Napkin Math 12.3: Payback of a debt cleanup

Problem: A team spends 40 hours/month manually fixing "broken plumbing" (stale data, failed scripts, manual monitoring). A one-month intensive cleanup (160 hours) is projected to reduce this to 8 hours/month. Is the cleanup worth it over 3 years of model lifecycle?

Math:

1. **Status quo** (3 years): 36 months × 40 hours/month × $150/hr = $216000.
2. **Proactive Path**: (160 hours investment + 36 months × 8 hours/month) × $150/hr = $67200.
3. **Net Savings**: $216000 - $67200 = $148800.
4. **Cost ratio**: 3.2×.

Systems insight: Under the projected toil reduction, proactive maintenance reduces the modeled cost by a factor of 3.2× over the model lifecycle. The benefit comes from avoiding repeated work long enough to recover the initial repair effort. A shorter remaining lifetime, lower recurrence, or additional upkeep can change the decision, so refactoring earns priority through measured recurring cost and a credible remediation plan.

The cleanup calculation becomes actionable only when the platform can compare unlike debts on a common operational scale. ML technical debt manifests in measurable symptoms that directly affect platform velocity and reliability (Sculley et al. 2015; Amershi et al. 2019).

12.2.4.1 Debt categories and measurement

The debt category matters because each failure mode leaves a different operational trace. Table 12.3 turns the taxonomy into a measurement map: data debt appears as incidents and manual pipeline intervention, configuration debt as unvalidated release surface, model debt as glue-code drag, and infrastructure debt as toil that consumes platform capacity.

Table 12.3: Technical Debt Measurement Map: Different debt types require different observable signals. Numerical thresholds are illustrative platform policies, not empirical laws. A shared scoring process can compare them only after each category has been tied to the symptom and threshold that make the debt operationally visible.

Debt type	Operational symptom	Measurement signal	Warning threshold
Data debt	Unstable dependencies, missing versioning, weak validation	Data incidents per month and manual pipeline intervention	More than 10 data incidents/month or 30% manual runs.
Configuration debt	Ad-hoc files, duplicated parameters, absent validation	Deployment failures, lines of config, unvalidated parameters	More than 500 lines of unvalidated configuration per model.
Model debt	Glue code, undeclared consumers, tangled serving paths	Coupling score, undocumented consumers, trace time	More than 20% engineering time maintaining glue code.
Infrastructure debt	Brittle pipelines, manual deployment, environment drift	Toil hours, automation coverage, drift incidents	More than 50% platform capacity spent on toil.

Four fleet-level indicators show when that debt is accumulating, whatever its category: deployment velocity from commit to production, incidents per thousand deployments, the share of team capacity spent on toil, and the share of dependencies behind supported versions. Each needs a baseline and an escalation threshold calibrated to the platform's workload risk and release cadence. The worked example starts from such symptoms and ranks the debts behind them with an ordinal score.

12.2.4.2 ML debt audit and prioritization

An ML platform team supporting 40 production models with 15 engineers faces deployment velocity problems. New models require 6 weeks to reach production, frustrating both platform and model teams. The audit estimates how much capacity each repair could release and compares that benefit with implementation effort and operational risk. A highly visible symptom is a starting point for measurement, not sufficient evidence of priority.

The audit applies the measurement map of table 12.3 to the deployment-velocity complaint. Before scoring, it records in table 12.4 the three categories that the team's release and incident symptoms point to: configuration debt, model debt in the form of pipeline glue code, and infrastructure debt in the form of fragmented monitoring. Data debt can also delay releases, but this audit has not attributed the observed delay to data incidents and leaves that category outside its scope. Each row pairs an operational symptom with the impact measure that makes the debt comparable across teams.

Table 12.4: ML Debt Audit Inputs: Initial observations for configuration debt, model debt (pipeline glue code), and infrastructure debt (monitoring) before priority scoring. The table preserves the symptom, impact metric, technical measure, and estimated annual cost for each debt category.

Debt category	Symptom	Impact metric	Technical measure	Estimated annual cost
Configuration debt	Each model has custom YAML configuration files averaging 847 lines with no validation schema	35% of deployment delays result from late config errors	Manual config review required for every deployment	12 engineer-hours per deployment $\times$ 80 deployments/year = 960 hours
Model debt (pipeline glue code)	Data preprocessing uses 23 different scripts with 62% code duplication	12 engineer-hours per week debugging pipeline breaks	No shared preprocessing library; each team implements custom logic	12 hours/week $\times$ 52 weeks = 624 hours

Debt category	Symptom	Impact metric	Technical measure	Estimated annual cost
Infrastructure debt (monitoring)	Each model uses ad-hoc monitoring, with no unified observability platform	Mean time to detect incidents is 4.2 hours	23 monitoring approaches across 40 models	Extended incident duration costs \$50K per incident × 15 incidents/year = \$750K

The scoring approach uses three criteria: impact severity, frequency, and resolution cost, each scored 1 to 3. Higher impact and frequency raise priority; higher resolution cost lowers it. Applied to the three debt categories observed earlier, these scores produce the ranking in table 12.5.

Table 12.5: Debt Prioritization Scoring: Three-criterion scoring (impact severity, frequency, and resolution cost) applied to three observed debt categories. Scores use Impact × Frequency/Resolution Cost, so slower or more expensive fixes are penalized. Configuration debt scores highest, justifying first-priority remediation.

Debt category	Impact	Frequency	Resolution cost	Total score	Priority
Configuration	3 (High)	3 (Daily)	2 (Medium: 6 weeks)	4.5	**1st**
Monitoring	3 (High)	2 (Weekly)	2 (Medium: 8 weeks)	3	**2nd**
Pipeline Glue	2 (Medium)	2 (Weekly)	3 (High: 16 weeks)	1.3	**3rd**

The score makes configuration debt the first paydown target: build configuration schema validation and a templating system before addressing lower-scoring pipeline glue. The scores are ordinal judgments rather than dollar estimates, so the larger estimated annual cost of monitoring debt in table 12.4 does not by itself raise its rank. Configuration debt ranks first because it recurs more often than monitoring debt at the same ordinal resolution cost. The investment requires one engineer working for 6 weeks to build the configuration system. The savings estimate needs one more input, the share of each deployment's manual configuration review that schema validation eliminates. Take it as 35 percent. That planning assumption happens to equal the share of deployment delays the audit attributes to configuration, but delays and review hours are different quantities, and only a pilot of the schema system can measure the second. With that input, released capacity is: 12 engineer-hours per deployment × 80 deployments per year × 35 percent = 336 engineer-hours saved per year. At \$150/hour, that is \$50K in annual savings, so the investment pays back in about 8.6 months.

12.2.4.3 Decision framework

The table score generalizes by adding a Benefit term that rates what the repair returns, such as released hours or avoided incidents. The result remains a ranking heuristic rather than a monetary return. Impact and Frequency are ordinal weights, Benefit is a separate score, and Resolution Cost is effort on a consistent internal scale. Because Frequency already appears as its own factor, Benefit must be scored per occurrence; an annual dollar benefit already includes frequency and would count it twice:

$$\text{Paydown Priority} = \frac{\text{Impact} \times \text{Frequency} \times \text{Benefit}}{\text{Resolution Cost}} \tag{12.2}$$

The paydown-priority formula in equation 12.2 keeps automation work tied to operational value: high-impact, frequent, high-benefit fixes outrank expensive work whose payoff is speculative. The score only ranks candidates. The payback calculation in the deployment-automation notebook below is a separate tool that keeps the terms with measurable units (deployments per week, toil hours saved per deployment, and setup hours). Any severity that saved hours cannot express, such as outage or compliance exposure, must therefore be argued alongside the payback period rather than folded into it.

Napkin Math 12.4: ROI of automation

Problem: A team spends 10 hours of manual toil per model deployment. Investing 120 hours in a CI/CD pipeline is projected to reduce deployment toil to 0.5 hours. At 3 deploys/week, how long until the automation pays for itself?

Math:

1. **Automation Cost**: 120 hours $\times$ \$150/hr = \$18000.
2. **Weekly Savings**: 9.5 hours/deploy $\times$ 3 deploys/week $\times$ \$150/hr = \$4275/week.
3. **Payback Period**: \$18000 / \$4275/week $\approx$ 4.2 weeks.

Systems insight: The projected payback period of 4.2 weeks makes this recurring deployment task a strong automation candidate. The calculation assumes that deployment frequency and time savings persist, and it excludes additional maintenance of the automation. Comparing that upkeep with the recurring effort avoided determines whether the initial payback becomes a lasting reduction in toil.

The score ranks candidates for investigation; comparing debt paydown with feature development then requires costs, expected benefits, uncertainty, and time horizons in common units. The configuration example scores highly under its stated assumptions; it is not automatically preferable to an unrelated project with a different scoring scale. A planned system retirement or a small affected population can reduce its practical value.

The same prioritization logic must become organizational habit, or the measured debt will reaccumulate between audits. Quantified debt enters a backlog with its affected systems, impact, resolution cost, and score, and a quarterly review re-ranks it against current deployment velocity, incident rates, and toil. A debt budget reserves an explicit, revisable share of engineering time for remediation, and design review applies the same reasoning upstream by catching new configuration complexity and hard-to-maintain dependencies before they become recurring toil.

12.2.5 When to form a dedicated platform team

Auditing and paying down cross-model debt eventually requires an owner with a charter wider than any single product team, because product teams rarely fund shared infrastructure on their own until an outage exposes undocumented dependencies, missing rollback paths, and idle GPU silos. A dedicated ML platform team is justified when fragmentation costs more than the shared infrastructure needed to replace it, combining quantitative factors (cost savings, velocity improvements) with qualitative factors (consistency, governance, talent retention).

The savings of shared capacity and the benefit-cost ratio point to three potential benefits of shared ownership: infrastructure efficiency, faster releases, and fewer incidents. Only the first is computed so far, in notebook 12.1, and it holds only while peak demand and reliability headroom permit the smaller pool. Faster delivery and fewer incidents are additional hypotheses to measure; neither a guaranteed engagement gain nor a fixed incident-reduction percentage follows from introducing a platform. The business case combines observed capacity savings with independently valued release and reliability improvements.

The qualitative case follows from the same fragmentation cost. Shared ownership gives telemetry, rollback capability, and runbook documentation a common acceptance process, but those standards protect the fleet only when deployment gates enforce them. A centralized platform team also gives product teams access to operational expertise and provides dedicated roles for engineers focused on machine learning systems and reliability. Platform-level lineage and policy controls support governance review by preserving evidence of which versions ran, what checks passed, and who authorized the release.

While the economic justification for platform operations becomes clear at scale, the technical implementation begins with a harder problem, preventing independent models from colliding through shared state. Multi-model management requires untangling the hidden dependencies that emerge when hundreds of models share the same data, infrastructure, and user experiences.

12.3 Governance of State and Change

Collisions through shared state rarely announce themselves at the model that changed. Imagine an e-commerce platform where the search ranking model uses outputs from a user embedding model: an unannounced dimensionality change can break the consumer's input interface, while a shift in scale or meaning can pass shape checks yet degrade ranking quality. Because shared features, chained predictions, and common infrastructure make model behaviors interdependent across team boundaries, governing a portfolio requires both an artifact record and a data contract. Registries, dependency tracking, and ensemble-aware lifecycle controls identify the model versions and consumers involved, while feature operations preserve the meaning and timing of the inputs those consumers expect so release gates can evaluate a complete dependency state rather than an isolated model binary.

12.3.1 Model registries at scale

Storing model weights in ad-hoc cloud buckets or team-local folders leaves operators unable to tell which commit or dataset produced a deployed checkpoint, or which downstream rankers will break when an upstream embedding model is updated. Turning checkpoints into governed interfaces requires a model registry that serves as the central catalog for all machine learning artifacts in an organization, and at enterprise scale it also exposes the dependency graph that determines which downstream systems an update can break.

12.3.1.1 Core registry requirements

The registry can enforce dependency control only if four ordinary catalog capabilities are reliable (table 12.6). Version identifiers let downstream consumers pin the exact artifact they validated; metadata exposes the training and deployment context needed for review; artifact storage turns registry entries into durable deployable assets; and access control keeps ownership boundaries explicit.

Table 12.6: Core Registry Requirements: Enterprise model registries need versioning, metadata, artifact storage, and access control before dependency-aware validation can be reliable. These catalog capabilities make model artifacts behave like governed production interfaces rather than files on shared storage.

Requirement	What it stores or enforces	Why it matters at scale
Version management	Unique artifact identifiers plus lineage for training data, hyperparameters, code commits, and evaluation metrics.	Downstream consumers can pin the exact artifact they validated.
Metadata storage	Training configuration, evaluation results, hardware requirements, serving configuration, and ownership.	Deployment and review decisions can be made without reverse-engineering the artifact.
Artifact storage	Durable model binaries with efficient retrieval and caching near serving locations.	Large models such as LLMs can exceed 100 GB and cannot rely on ad-hoc file distribution.
Access control	Read-write, administrative, and read-only permissions at team and dependency boundaries.	Model developers, platform operators, and dependent teams interact through governed APIs.

12.3.1.2 Dependency tracking

Beyond these core requirements, the distinguishing feature of enterprise registries is explicit dependency tracking. When Model A consumes features computed by Model B, that relationship must be recorded and enforced, because the recorded edges are what bound the work an upstream change creates. A version bump on embedding model E propagates along those edges to an impact set of candidate retrieval, parallel rankers, and an ensemble, all of which must be re-evaluated before release (figure 12.3).

The impact boundary in figure 12.3 identifies consumers that require evaluation after an upstream change. Fraud model F has no recorded dependency, so the traversal does not flag it; an omitted edge would leave that conclusion incomplete. Within the recorded graph, upstream embedding model E produces representations consumed by retrieval model R, three parallel rankers (K_1, K_2,

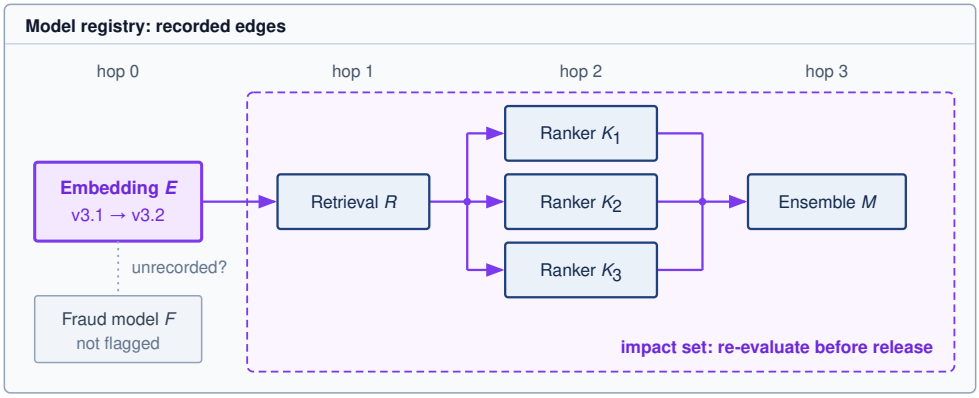

Figure 12.3: Dependency-Aware Model Registry: A version bump on embedding model E starts a traversal of the registry's recorded downstream edges: retrieval R, rankers K_1 to K_3, and ensemble M form the impact set (dashed region) that must be re-evaluated before E v3.2 ships. A model with no recorded edge from E, such as fraud model F, is not flagged, which is why an unrecorded dependency is a silent risk. The rankers also consume features derived from E directly; only the retrieval path is drawn.

K_3) score those candidates using features derived from E, and ensemble model M combines the scores into a final ranking. Recording these relationships tells the release controller where to test, and the tests on those consumers decide whether the new representations remain compatible.

Update approval begins with graph traversal and then requires behavioral and interface evidence. When embedding model E is updated, the registry and release controller coordinate four steps:

1. Traverse downstream edges to identify every dependent model (R, K_1, K_2, K_3, and M).
2. Re-evaluate each dependent model against the proposed embeddings.
3. Block deployment of E until every consumer passes.
4. Roll out the approved version set through compatible intermediate states, retaining a validated rollback target. Topological ordering, which updates every producer before its consumers, is insufficient on its own if an existing consumer cannot accept the new producer's outputs.

Without explicit dependency tracking, organizations discover architectural coupling through production outages when an upstream change silently corrupts downstream assumptions. A single YAML entry captures artifact location, training provenance, evaluation results, and declared dependencies (listing 12.1). The downstream entries seed impact analysis; following their registry records reveals transitive consumers. Release approval additionally requires confirming that the update preserves the input-output contract expected by those consumers, or scheduling a coordinated rollout while retaining a validated rollback target.

Listing 12.1: Model Registry Schema: A YAML entry capturing model metadata, artifact location, training provenance, and evaluation results for dependency tracking and reproducibility.

```yaml
model:
  name: user_embedding_v3
  version: "3.2.1"
  type: embedding_model
  domain: recommendation

artifact:
  path: gs://models/user_embedding_v3/3.2.1/
  format: tensorflow_savedmodel
  size_bytes: 4294967296

training:
  data_version: user_interaction_2024_01
  code_commit: abc123def
  started_at: 2024-01-15T10:00:00Z
  duration_hours: 48
  hardware: 8xA100-80 GB

evaluation:
  metrics:
    recall_at_100: 0.342
    embedding_quality: 0.891
  evaluation_set: eval_2024_01

dependencies:
  upstream:
    - feature_store/user_features_v2
    - feature_store/interaction_features_v1
  downstream:
    - models/candidate_retrieval_v4
    - models/ranking_ensemble_v2

serving:
  min_replicas: 10
  max_replicas: 100
  latency_p99_target_ms: 5
  memory_gb: 16

ownership:
  team: recommendation-core
  oncall: recsys-oncall-team
```

12.3.2 Ensemble management

The downstream entries in listing 12.1, candidate retrieval and a ranking ensemble, point to a composition problem that registry edges alone cannot solve. When a recommendation pipeline chains retrieval, ranking, and filtering models owned by different teams, an upstream retrieval change that improves local recall can shift the candidate distribution enough to degrade final ranking. Recommendation systems exemplify this multi-model management challenge by chaining stages in sequence and combining predictions from several models within a stage. The platform must validate the complete recommendation interface even when different teams own its components.

12.3.2.1 Architectural rationale for recommendation ensembles

Recommendation architectures separate work when evaluating every candidate with the richest model would exceed the latency or compute budget. Retrieval produces a manageable candidate set, ranking scores that set with richer features, and later re-ranking or business-rule stages produce the final ordering (Covington et al. 2016; Liu et al. 2022). Within this path, specialized models can address competing objectives such as engagement, diversity, and freshness, while ensemble stages combine their outputs. The decomposition trades a smaller amount of expensive scoring for additional interfaces whose behavior the platform must validate.

The same decomposition improves experimentation velocity because teams can update individual components without retraining the entire recommendation stack, but that independence creates version-compatibility obligations for the platform. It also changes risk management. Other ensemble components can sometimes compensate when one model fails or produces poor results, yet the platform must detect when compensation hides a regression rather than resolving it.

12.3.2.2 Ensemble deployment patterns

Deploying an ensemble update requires evaluating its effect on the composed system. Consider updating the fine ranking model within a recommendation ensemble. An illustrative staged ensemble deployment moves the update through four phases (table 12.7): shadow deployment (24–48 hours logging without serving), canary (1 percent traffic for 4–8 hours), staged rollout (5 percent to 100 percent over 24–72 hours), and soak (7–14 days monitoring for delayed effects). These durations are planning assumptions; promotion requires enough eligible observations, mature outcomes, and acceptable guardrail metrics for the release decision.

Table 12.7: Staged Ensemble Deployment Pattern: Four-phase rollout for updating recommendation ensemble components. Shadow deployment (24–48 hours) compares behavior without serving candidate outputs; resource isolation and outcome evidence bound what it establishes. Canary (1 percent traffic, 4–8 hours) enables statistical regression detection, staged rollout progressively increases exposure, and soak period (7–14 days) catches delayed interaction effects.

Deployment stage	Actions	Duration	Rollback trigger
Shadow	New model scores alongside production, results logged but not served	24–48 hours	Quality metrics below threshold
Canary	1% traffic receives new model results	4–8 hours	Statistical significance of regression
Staged Rollout	5% → 25% → 50% → 100%	24–72 hours	Business metric degradation
Soak	Full traffic, extended monitoring	7–14 days	Delayed effects emerge

The extended timeline reflects the difficulty of detecting regressions in ensemble systems. A component change that improves its local metrics might degrade system-level performance through subtle interactions with other components.

12.3.2.3 Interaction effects

Ensemble components interact in ways that make local validation insufficient. If an upstream retrieval bug weakens the semantic similarity score passed to the ranking stage, a ranking model retrained during that period learns to down-weight the noisy retrieval score and rely more heavily on auxiliary popularity or recency signals. When engineers later fix the retrieval bug and restore sharp candidate sets, the adapted ranker still underweights semantic relevance, so an upstream improvement can paradoxically degrade end-to-end ranking quality. Even without retraining, a fixed ranker can respond poorly when an upstream update alters the feature distributions or candidate mix it was calibrated to score. Over longer horizons, ranking decisions also change which items users encounter, shaping the click and conversion feedback collected for future training. Logging intermediate candidate sets and stage scores helps isolate where these shifts enter the pipeline, while stable holdouts (fixed user groups kept on the previous system for comparison) and extended observation windows evaluate the complete system. Joint retraining is a targeted intervention when diagnostics confirm that downstream weights have adapted to an outdated upstream distribution, rather than a routine periodic reset.

12.3.3 Model lifecycle management

Governing those dependencies over time requires controls at both ends of a model's life. Promoting experimental notebooks straight to live traffic without staging gates causes avoidable outages. Retiring a model without migrating its consumers leaves behind "zombie models"[2] that silently drain GPU capacity and block platform upgrades. Lifecycle management keeps model dependencies from becoming permanent obligations: every promotion widens a model's blast radius, and every

[2] **Zombie models**: Production models that continue serving despite being obsolete or superseded. Even a low-traffic model consumes registry, monitoring, dependency-management, and on-call capacity until it is deliberately retired. The operational cost therefore extends beyond compute: zombie models enlarge the dependency graph, making platform-wide upgrades and security patches slower and riskier.

retirement must migrate its consumers before the platform can reclaim resources. The lifecycle stages trace that progression from development to archive, each defined by the gates it imposes rather than by a status label.

```
Development → Staging → Canary → Production → Deprecation → Archive
```

In development, models exist as experimental artifacts. Operations requirements are intentionally light, but they are not optional: successful experiments must preserve results, version history, and enough reproducibility metadata to be challenged later. The operational concern is therefore transition readiness. A promising model should not move to staging until it has clear production-readiness criteria, automated evaluation against production-equivalent data, and documentation sufficient for review.

Staging tests the candidate as a deployable service before it serves user-visible outputs. Production-equivalent data, feature interfaces, and hardware make latency, throughput, and behavioral checks relevant to deployment. Shadow traffic can add evidence about live request patterns when its coverage justifies its cost. The promotion gate combines the checks required by the release contract with human review where model behavior or risk demands judgment.

Production widens the blast radius from validation traffic to live traffic, first through a small canary slice (section 12.5.5.2) and then at full exposure, so the model now requires continuous monitoring with alerting, capacity for traffic fluctuations, rollback procedures, and on-call support. Production is not a terminal state. Distribution shifts prompt investigation and evaluation; retraining is warranted when the resulting evidence supports an update. Feature and infrastructure changes likewise require compatibility checks, and periodic comparison with newer baselines tests whether the deployed model still earns its place.

The lifecycle closes only when a model can be retired without breaking its consumers. Deprecation identifies dependent systems, provides migration paths and timelines, maintains the old model until migration completes, and archives artifacts for reproducibility and audit. Organizations often underinvest in this final stage, allowing obsolete models to linger in production long after their replacements ship. Platform-level lifecycle enforcement closes that gap.

12.3.4 Deployment patterns by model count

How a model moves through those lifecycle stages depends on how many other models move with it. Rolling out a single isolated vision classifier once a month requires only a local canary, whereas pushing daily updates across a 50-model recommendation ensemble or hundreds of platform models without cross-model rate limits can trigger cascading regressions. The appropriate deployment pattern depends on the interdependence of the versions being updated, because a compatibility change can widen the rollback unit. Across the example deployment patterns in table 12.8, single model deployments update monthly for isolated vision classifiers, pipeline deployments weekly for 3–5 sequential NLP models, ensemble deployments daily for 10–50 recommendation components, and platform deployments continuously across hundreds of enterprise models. The cadences describe example operating contexts; model count does not determine release frequency or the evidence required for promotion.

Table 12.8: Deployment Patterns by Scale: Four patterns addressing different model counts and update frequencies. Single model deployments (1 model, monthly updates) use standard canary rollouts, while platform deployments (100+ models, continuous updates) require automated policy enforcement, cross-model impact analysis, and global rate limiting to prevent simultaneous high-risk deployments.

Pattern	Model count	Update frequency	Example
Single Model	1	Monthly	Vision classifier
Pipeline	3-5	Weekly	NLP processing pipeline
Ensemble	10-50	Daily	Recommendation system
Platform	100s	Continuous	Enterprise ML platform

For a model without downstream model consumers, the rollback unit can be smaller. Canary deployments, blue-green switches,[3] and gradual rollouts still require a validated previous version

[3] **Blue-green deployment**: A release pattern that prepares the new version in a parallel environment before switching production traffic. Fast rollback requires the previous environment to remain ready and its dependencies to remain compatible. Holding both versions during the transition adds serving capacity and cost (C_{infer}); the duplication depends on which resources the environments can safely share.

together with compatible preprocessing, features, and serving configuration. Restoring the model artifact alone is sufficient only when those surrounding interfaces remain compatible.

Pipelines introduce ordering constraints because models execute in sequence and each model's output feeds the next. Compatibility must hold at every intermediate deployment state. An upstream-first update is safe only if existing consumers accept its new outputs; otherwise the platform must prepare compatible consumers or stage both versions before switching traffic. The release record names the compatible version set and the rollback unit, which may span the entire pipeline.

Ensembles widen the coordination problem because multiple models may execute in parallel or in a graph. Different teams can update components on different schedules, partial updates may change only part of the ensemble, and the system behavior emerges from component interactions. Testing in isolation is therefore insufficient; integration testing becomes the deployment gate.

At platform scale, continuous deployment means some model is always being updated somewhere, so the platform must prevent individually safe changes from combining into an unsafe fleet state. Platform deployment requires automated rollout policies based on model risk classification, cross-model impact analysis before deployment approval, global rate limiting to prevent simultaneous high-risk deployments, and automated correlation of incidents with recent deployments.

12.3.5 E-commerce model ecosystem

The registry, ensemble, lifecycle, and deployment-tier controls developed across this section converge in the e-commerce dependency graph that opened it, where updating a single user-embedding model alters the inputs to candidate retrieval, price sensitivity, ranking, and business rules simultaneously and turns a local model upgrade into a five-tier integration test. An e-commerce platform typically starts the graph with embedding models. A user embedding model generates user representations from behavior history, while a product embedding model represents products from attributes and interactions. Those embeddings are not final predictions; they are shared interfaces consumed by downstream services.

The next layer turns shared representations into candidate sets and business signals. A candidate retrieval model uses the embeddings to find relevant products, and a price sensitivity model estimates how strongly a user will respond to price changes. A ranking model then scores the candidate products using both embedding features and auxiliary signals, while a business rules model applies promotional and inventory constraints before results reach the user.

The resulting graph has five tiers, and its shape decides what a change costs. User and product representations travel over a shared bus into candidate retrieval and price sensitivity, which converge in ranking before business rules apply (figure 12.4).

Because the embedding bus in figure 12.4 branches directly to both mid-tier consumers, a change can affect candidate retrieval and price sensitivity before propagating through ranking and business rules. The platform therefore evaluates both direct embedding compatibility and downstream recommendation quality. If old consumers cannot accept the new embeddings, the release can retain both versions and route each consumer to the version it validated while migration proceeds. The rollback plan must preserve that compatible state as well. Because these models share features in addition to embeddings, the release must also pin versioned feature definitions and carry their guarantees of point-in-time correctness and freshness, the contracts that complete the state a release pipeline must validate.

12.4 Feature Store Operations

Governing model checkpoints and their dependency edges still leaves the input values those models ingest unprotected if training and serving compute features through separate code paths. Consider a fraud detection system that needs to know a user's transaction volume over the last ten minutes. Serving retrieves that feature within the prediction's latency budget, while training reconstructs it for many historical transactions. These access paths differ, but the feature's meaning must remain compatible. If training counts transactions that were unavailable at the historical prediction time, or serving applies a different aggregation rule, the model receives inputs that differ from those on which it was evaluated.

A common design separates the two access paths: online stores serve low-latency feature lookups, while offline stores generate reproducible training data. The two paths form a training-serving gap

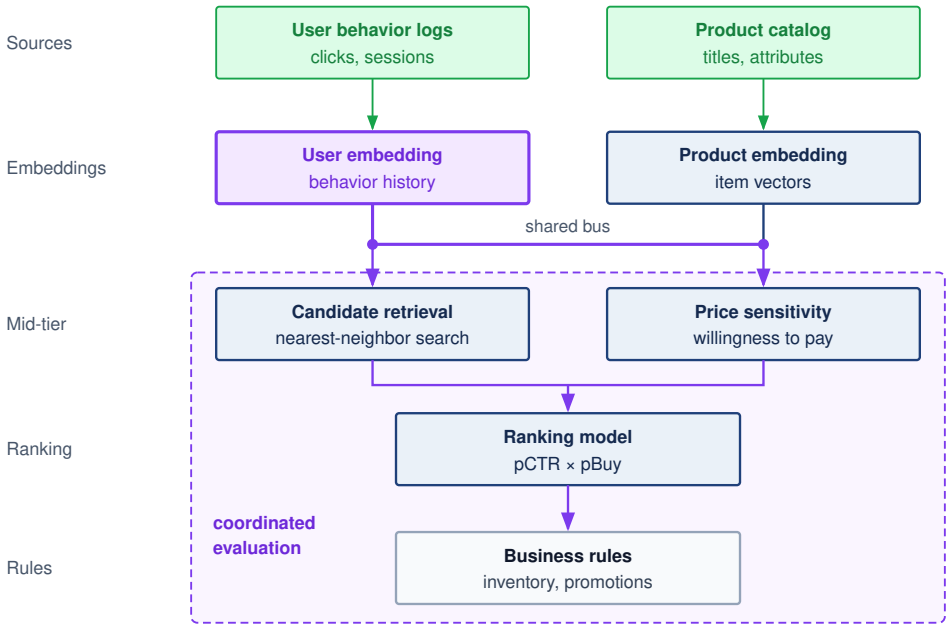

Figure 12.4: E-Commerce Model Ecosystem: Data sources feed a user embedding and a product embedding. A shared bus carries both embeddings into both mid-tier models, candidate retrieval and price sensitivity, whose outputs merge in the ranking model before business rules apply inventory and promotion constraints. Because every embedding reaches every downstream model, a change to the user embedding requires coordinated evaluation of all four consumers.

in computational context. Training computes features in *batch* over *hours*, while serving computes them in *real time* within *milliseconds*, yet each feature must mean the same thing on both paths. The gap is a design condition, not yet a failure. It becomes one as training-serving skew, the failure mode in which small differences in processing logic between the two paths give the deployed model inputs it never saw in training and degrade accuracy silently.

> **Systems Perspective 12.2: Training-serving skew**
>
> At fleet scale, the concern shifts from detecting training-serving skew in a single model to controlling it as a platform-level property. A feature store reduces skew by making training and serving pipelines read from shared feature definitions and materialized feature values where possible, turning $f_{\text{train}}(x) \equiv f_{\text{serve}}(x)$ into an architectural invariant that can be tested and monitored rather than a manual convention repeated by every team.

12.4.1 Feature store architecture

Training jobs scan terabytes of historical features across millions of examples, whereas live serving must fetch a single user's latest feature vector in under 10 milliseconds without diverging in feature logic. The **dual-store architecture** resolves this conflict between high-throughput batch scans and low-latency point lookups by splitting storage while the platform enforces compatible definitions and temporal semantics across both paths. One materialization engine, governed by versioned feature definitions, writes ingested batch and stream data into an offline analytical store for historical joins and an online key-value store for point lookups (figure 12.5).

The **offline store** retains historical feature data for reproducible training examples, so its access path favors large scans and joins, where a single query might read terabytes to build a feature set for

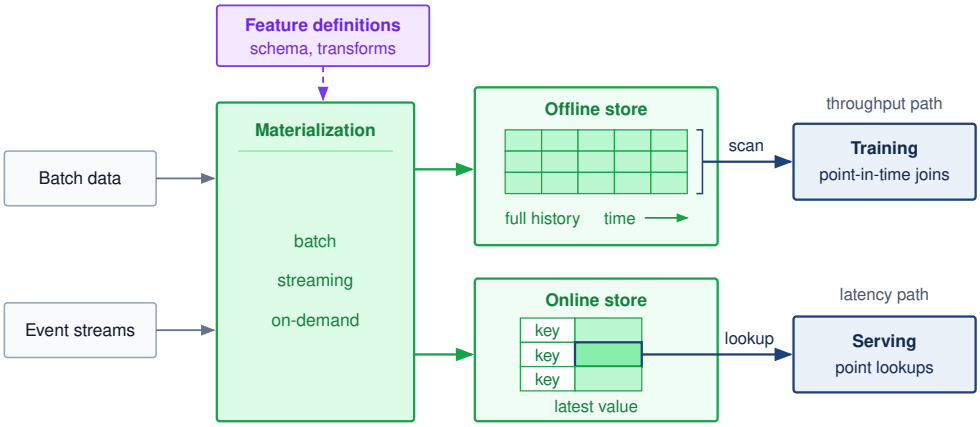

Figure 12.5: Feature Store Architecture: A dual-store architecture resolves the conflict between training (high-throughput batch scans) and serving (low-latency point lookups). One set of feature definitions drives materialization from batch and streaming sources into an offline store, which keeps full history for point-in-time training joins (the throughput path, top), and an online store, which keeps the latest value per entity for serving lookups (the latency path, bottom). Shared definitions support compatible feature meanings; materialization and historical-availability checks must verify that contract.

millions of examples. Throughput determines how quickly that dataset can be prepared within the training schedule.

The **online store** serves feature values within the prediction's latency budget. At fleet scale its load depends on request rate, features fetched per request, batching, and cache reuse, rather than model count alone, and the storage choice must satisfy that access pattern at the required tail latency and peak load.

Materialization loads computed features into both stores: in batch on a daily or hourly schedule, from event streams in near real time, or on demand at request time when freshness requirements exceed batch frequency. The central architectural challenge is ensuring consistency between the two stores during materialization. If the offline store contains a feature value for training that was not identically available in the online store at the time of prediction, it creates a training-serving mismatch. That mismatch can silently degrade model performance in production, which is the training-serving skew problem formalized in perspective 12.2.

For a platform serving hundreds of models, such as the five-hundred-model portfolio that opened this chapter, shared feature contracts prevent each consumer from reconstructing historical correctness and serving consistency independently. The dual-store pattern provides one implementation of that contract; the invariant is compatible feature meaning and temporal availability across both paths, rather than a requirement that every workload use the same storage system.

12.4.2 Freshness SLOs

Keeping the online store aligned with recent user behavior requires specifying how quickly new events must cross that materialization path. Recommending products from a 24-hour batch feature table ignores everything a shopper clicked in the last ten minutes, whereas streaming every static demographic attribute at one-second latency wastes compute on values that rarely change. Feature freshness represents the delay between real-world events and their reflection in feature values. Example requirements range from days for static attributes such as demographics, which batch computation can serve, to seconds for features built on the last user action, which may require streaming or on-demand computation (table 12.9). The consuming decision determines the acceptable delay; the feature category alone does not establish a service-level objective (SLO).[4] The categories in the table are a starting point for those SLOs; monitoring and alerting must use the timestamp and delay contract validated for the consuming workload.

[4] **SLO**: A target for a declared service indicator and measurement window. Under a time-based definition, "three nines" (99.9 percent) annual availability allows about 8.8 hours of downtime; a request-based definition instead counts failed eligible requests. A p99 latency target of 200 ms requires at least 99 percent of eligible requests to complete within that bound over the specified window.

Table 12.9: Feature Freshness Requirements by Type: Four feature categories with SLO thresholds and computation patterns. Static features (user demographics) tolerate day-scale staleness with batch computation; real-time features (last user action) demand seconds-scale freshness through streaming or on-demand computation, directly impacting recommendation quality and engagement.

Feature type	Example	Freshness SLO	Computation pattern
Static	User demographics	Days	Batch
Slowly changing	User preferences	Hours	Batch
Session-level	Current session context	Minutes	Streaming
Real-time	Last action	Seconds	Streaming/On-demand

The table's freshness classes are budgets on feature freshness latency, the event-to-availability delay between a user action and its arrival in the model's inputs. Stale features can degrade relevance even when serving remains fast and available.

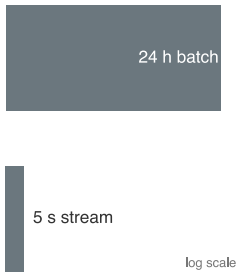

Streaming closes the freshness lag that batch leaves open.

Example 12.1: Feature freshness latency

Scenario: A real-time recommendation service processes user interaction logs to compute feature vectors during an active user session.

Diagnosis: Suppose the daily batch pipeline delivers relevant interactions with freshness latency $T_{\text{freshness}} \approx 12$–$24$ h. Those features cannot reflect activity from the current session, even if other recommendation inputs change. A streaming pipeline (Kafka/Flink) operating at $T_{\text{freshness}} \approx 1$–$5$ s can incorporate that activity sooner; whether it updates the vector before the next user action depends on the session's timing.

Systems lesson: Feature freshness latency is a physical system metric ($T_{\text{freshness}} = t - T_{\text{event}}$). Reducing freshness latency can improve decisions that depend on recent behavior, but engagement lift must be measured in an experiment. The business benefit must justify the additional streaming capacity and operational complexity.

Recommendation systems show how a missed freshness SLO reaches the user when a ranking decision depends on behavior from the current session. At platform scale, that session-level deadline must hold across multi-terabyte embedding tables and millions of queries per second while sharing ingestion capacity with offline training.

Lighthouse 12.2: Archetype B (DLRM at Scale): Feature staleness

Context: Archetype B (DLRM at Scale) recommendation workloads maintain sparse embedding tables of 10 TB+ across distributed memory and serve on the order of 10 million queries per second platform-wide, each within a p99 serving budget ($L_{\text{lat,SLO}}$) of 100 ms, of which feature retrieval receives only a few milliseconds.

Mechanism: Ground truth user intent shifts continuously. If the feature materialization pipeline adds a 10-minute update lag ($T_{\text{staleness}} = 600$ s) to freshness latency, serving queries read feature values that omit the last 10 minutes of user behavior, so downstream rankers can recommend items the user has already moved past.

Impact: The lag is dominated by pipeline scheduling rather than transfer time, and it can make recent behavior unavailable to a ranking decision. The resulting click-through or revenue effect must be measured for the service; there is no universal staleness penalty.

Systems lesson: When recent behavior materially affects ranking quality, DLRM operations must provision ingestion and materialization to meet the feature's freshness SLO. Scheduling must preserve that deadline alongside offline training needs. The measured value of fresher inputs determines whether the additional streaming capacity earns its cost.

Freshness becomes operational only when the SLO names its timestamp contract. Equation 12.3 measures refresh age, the time since the most recent feature update, which is a different quantity from freshness latency. Refresh age can stand in as a proxy for freshness latency against the thresholds in table 12.9 when updates certify that the source stream has been processed through the required time:

$$\text{Staleness} = t_{\text{current}} - t_{\text{feat}} \tag{12.3}$$

The two quantities diverge because rewriting an old source value can make its update timestamp recent without making the source information fresh. Streaming monitors therefore need source progress or watermarks as well as update timestamps, and must distinguish an idle entity from a stalled pipeline. Batch age increases between scheduled refreshes. Each path needs a threshold matched to its update contract rather than one global freshness rule.

12.4.2.1 Freshness impact on model quality

A recommendation system uses user interaction features with different freshness levels. Assume an online experiment yields the relative engagement lifts in table 12.10. These constructed values illustrate a planning calculation; replay of historical inputs alone does not establish a causal engagement gain:

Table 12.10: Assumed Engagement Lift by Feature Freshness: Constructed relative lifts measured against the same baseline engagement rate. The hourly-to-real-time gap is 2.1 percentage points of the relative-lift metric, not 2.1 percentage points of the raw engagement rate. The example translates that assumed gain into an infrastructure budget.

Feature freshness	Engagement lift vs. baseline
Real-time (< 1 min)	+12.3%
Near real-time (< 5 min)	+11.8%
Hourly	+10.2%
Daily	+8.1%

The difference in relative lift between hourly and real-time features is 2.1 percentage points of the shared baseline. If the organization values that incremental engagement at $10 million annually, the comparison must use the incremental cost of replacing the hourly path, including migration and ongoing operation over the same horizon. The engagement experiment supplies the quality difference; the business valuation and infrastructure cost determine whether it justifies investment.

Checkpoint 12.2: Pricing the streaming and batch freshness decision

Under the constructed lifts in table 12.10, moving from hourly to real-time features adds 2.1 percentage points of engagement lift. Treat that lift as $10 million in annual engagement value and apply it to the choice between streaming and batch infrastructure.

- ☐ **Freshness break-even**: determine the maximum annual incremental cost of a real-time feature pipeline relative to the hourly path that still breaks even on the 2.1-point lift, holding the $10 million engagement value constant, and state the cost threshold under which streaming is a clear win.
- ☐ **Lift per point**: compare the additional engagement lift for the daily-to-hourly, hourly-to-near-real-time, and near-real-time-to-real-time transitions, and identify the incremental pipeline-cost estimates needed to choose among them.
- ☐ **Scaling to users**: if the $10 million annual engagement value scales linearly with platform user count, determine the break-even user count as a fraction of the current population for a streaming feature platform costing an additional $5 million per year relative to the hourly path. State the range where it does not pay off, and choose an alternative architecture from table 12.9.

12.4.3 Point-in-time correctness

Joining a noon training example against a midnight batch table leaks six afternoon clicks into the model's training inputs, inflating `total_clicks_today` 2.5× and producing offline metrics that collapse in production. Training data must reconstruct features strictly as they existed at the timestamp of each historical event (figure 12.6). Using unaligned feature values as inputs for historical training examples creates data leakage[5] that inflates offline evaluation while leaving the deployed model without the future signals it learned to rely on.

Figure 12.6: Point-in-Time Correctness: A training example at noon should see only the clicks already logged by noon (solid green, 4), not the clicks that arrive later (dashed gray). A naive join that reads the midnight batch value of `total_clicks_today` counts all 10 and leaks six future events into the example. A point-in-time join pairs each training event with feature values as they existed *at the event timestamp*, so the example contains only information that would have been available at inference time.

Comparing the aggregation brackets in figure 12.6 isolates the leakage mechanism. The point-in-time join uses the four clicks logged before noon, whereas the naive join reads the midnight total, inflating `total_clicks_today` by 2.5×. This "time travel" failure uses a value that is valid in the database but invalid for the prediction moment: `total_clicks_today` is 10 at midnight, while the noon prediction can use only the 4 clicks already available. Training on the later value exposes information unavailable to the deployed predictor and can make offline evaluation overestimate production quality. Feature stores prevent that leakage by implementing **point-in-time joins** that retrieve feature values as of specific timestamps, as shown in listing 12.2.

Listing 12.2: Point-in-Time Join: A SQL lateral join that retrieves the most recent feature values available before each event, preventing future data leakage into training examples.

```
SELECT
    e.user_id,
    e.event_timestamp,
    e.label,
    f.feature_1,
    f.feature_2
FROM events e
LEFT JOIN LATERAL (
    SELECT feature_1, feature_2
    FROM features f
    WHERE f.user_id = e.user_id
      AND f.feature_timestamp <= e.event_timestamp
    ORDER BY f.feature_timestamp DESC
    LIMIT 1
) f ON TRUE
```

5 **Data leakage**: Future information can inflate offline metrics when a model learns signals unavailable at prediction time. Shared feature infrastructure should enforce historical availability boundaries, but model and data owners must still specify the prediction moment and verify how features were computed. A timestamped join cannot remove future information already embedded inside an incorrectly constructed feature.

The query assumes that `feature_timestamp` records when the value became available to serving. If it records only the source event time, a late arrival or later backfill could pass the filter despite being unavailable at the historical prediction. The store must retain availability time as well as event time when those differ, and the training join must enforce the availability boundary.

Point-in-time correctness has a storage cost because the store must retain feature history rather than current values alone. The same guarantee that prevents leakage also multiplies retained state:

$$\text{Storage} = N_{\text{entities}} \times N_{\text{features}} \times \frac{T_{\text{retention}}}{T_{\text{update}}} \times \text{bytes}_{\text{row}}$$

For 100 million users and 1,000 features, with a retention window of 1 year and an update interval of 1 hour, the history contains 876 trillion feature values. Each stored row carries more than the value itself. An 8-byte value needs an 8-byte entity key and, because a correct point-in-time join must filter on availability time as well as event time, two 8-byte timestamps, for 32 bytes per row. Full hourly snapshots then represent approximately 28 PB before compression, three quarters of it metadata that makes the history reproducible. Storing only changed values, compressing columns, and limiting retention can reduce this footprint. The reduction depends on change frequency and encoding, while retention must still cover the historical predictions the platform needs to reconstruct.

12.4.4 Feature versioning and lineage

Point-in-time joins keep a feature's historical values temporally consistent, but they cannot protect downstream consumers when a team rewrites a feature's logic in place. Suppose a platform team changes `user_engagement_score` from a 30-day z-score to a 7-day min-max scale. The output is still a floating-point number, so schema validation passes while the feature's meaning changes completely for every dependent model. Because feature definitions are production interfaces, versioning and lineage must make this evolution governable rather than implicit: versioning names the semantic change before it reaches production, lineage identifies every model trained or served with the old meaning, backfill recomputes historical values for the new meaning, and freshness and quality checks decide whether the rewritten history is safe to use. Without that chain, a feature migration becomes a silent multi-model regression.

A safe feature interface records both version dimensions and lineage fields, and each field supports a specific operation, such as blocking a silent formula change or reproducing a historical value for audit (table 12.11).

Table 12.11: Feature Version and Lineage Components: Feature versioning separates logic, data, and schema changes; lineage records the path needed to debug, audit, and reproduce feature values. Together, these fields turn a feature definition into a governed production interface.

Component	What it records	Operational use
Definition version	Computation logic for the feature	Blocks silent semantic changes when the formula changes but the type does not.
Data version	Source snapshot or stream contract	Distinguishes model changes from upstream data changes.
Schema version	Output type, shape, units, and nullability	Preserves API compatibility checks while still naming semantic shifts.
Consumer declaration	Exact feature versions each model consumes	Gives the platform a place to block unsafe updates before serving.
Source lineage	Source tables, streams, and their versions	Traces bad predictions back to the upstream data that produced them.
Transformation and runtime provenance	Transformation code, timestamp, and environment	Reproduces historical values for audit, debugging, and compliance.
Quality and freshness metrics	Validation results observed when the value was produced	Estimates blast radius and decides whether rewritten history is safe to use.

12.4.5 Backfill procedures

Once a new feature version is declared, dependent models cannot train on it until the platform populates its historical values. Recomputing a year of history without idempotent partition writes

or event-time bounds can saturate cluster compute or stamp training rows with post-hoc availability times. A backfill is a production migration over history. Changing a feature definition creates a new version of the training history for dependent models, so the platform must bound the compute cost, validate values against the intended definition, and preserve a path back. Treating the backfill as a migration changes the operational question from whether the computation can run once to whether history can be rewritten without corrupting future training data.

Feature replay requires a concrete execution safety contract. Each backfill partition needs a stable identity and an idempotent or transactional write path so retries cannot duplicate feature effects; event-time bounds and the late-data policy must also match the point-in-time join used for training. Otherwise, a mechanically successful replay can produce a history that is complete in processing time but incorrect at the prediction time the model is meant to reproduce.

At scale, historical data may have moved into cold storage, so recomputation first becomes a data-placement problem. The long historical window then competes with production workloads for compute. Even after the job finishes, the platform cannot yet trust the new values. It must compare overlapping periods against the intended new computation, explain differences from the old definition, and coordinate with dependent pipelines so training jobs do not mix old and new semantics mid-run.

↩ Exactly-once commit boundaries, event-time watermarks, and late-data policies for distributed transformations are established in section 4.7.1.

The practical pattern is incremental. Historical data is processed in bounded date partitions, and each partition is validated before consumers can use it. Independent partitions can run concurrently within the backfill's resource budget. During a dual-write period, old and new computations run in parallel, giving the platform overlapping values for comparison before cutover. The previous feature version and its lineage remain available through the migration and rollback window so dependent models can return to a compatible state.

12.4.6 Scale challenges

Once feature versions are backfilled and promoted, the online store must sustain live request volume within strict tail-latency bounds. Serving 1 billion daily recommendations that each fetch 100 features generates 100 billion feature lookups per day (roughly 1.16 million feature lookups per second before peak traffic), which is why feature stores at recommendation scale are critical-path serving infrastructure rather than a metadata convenience. This count measures feature-value accesses rather than backend requests; batching and cache reuse determine how many requests reach the store. A planning envelope of 5–10× the mean raises the required capacity accordingly; observed request patterns must establish the appropriate peak allowance.

Feature retrieval must fit inside the overall latency budget because lookup time on the critical path leaves less time for ranking. If the recommendation budget is 50 ms, the feature store may receive only 5–10 ms. Network overhead can consume 1–2 ms of that allocation, leaving roughly 3–8 ms for the store lookup itself. In-memory storage or caching, bounded batching, and placement near the serving fleet are candidate mechanisms. The selected design must meet the remaining lookup budget under peak load.

Production feature stores also carry two time horizons at once. Online capacity depends on the number of populated entity-feature pairs, encoded bytes per value, metadata, and replication. Offline capacity additionally depends on update frequency and retained history, as the snapshot calculation above shows. Neither footprint follows from model count alone. Multi-region replication can improve availability and lookup latency, but it also requires an explicit contract for how quickly new feature values become visible in each region.

12.4.7 Data quality operations

Meeting the millisecond lookup budget and multi-region replication contract still does not guarantee usable predictions if an upstream pipeline drops 18 percent of daily user events or leaks sentinel values like `-1` and `999` into an age feature while HTTP health checks stay green. Data quality issues are a major source of production ML problems (Polyzotis et al. 2017). Model metrics reveal downstream effects after users are exposed, whereas data checks can identify missing, stale, or invalid inputs earlier in the pipeline so validation gates can quarantine suspect data before consumers use it. At scale, these checks need explicit response policies so detecting an upstream problem leads to containment rather than merely another dashboard signal.

The first step is to quantify data quality in dimensions that support release decisions. Completeness measures whether expected records and fields are present. A daily data pipeline expected to produce 10 million user events but delivering only 8.2 million user events is 82 percent complete relative to that expectation. Investigation must distinguish missing delivery from a legitimate change in source volume. Consistency captures schema compliance and referential integrity; an age field with an allowed range of 0 to 120 would reject values such as 999 or -1, which may represent leaked sentinels. A policy might quarantine a batch when more than 5 percent of its checks fail, but a single critical violation can warrant rejection. The policy must weigh the consequences of invalid inputs against delayed or unavailable data.

Timeliness measures freshness against the SLO of the consuming model. Fraud detection may need features less than 100 milliseconds old, while demographic features can tolerate staleness measured in days; the alert threshold should therefore be tied to the feature's use, not a global constant. Accuracy asks whether values remain correct within expected distributions. Sensor drift after calibration lapses, for example, can leave values well-typed but wrong. Statistical tests such as Kolmogorov-Smirnov for continuous features, chi-square for categorical features, and maximum mean discrepancy for high-dimensional data detect distribution changes that warrant investigation; they do not establish that individual values are accurate.

Validation then turns those measurements into a layered defense. Schema validation enforces expected column names, types, and constraints at ingestion; tools such as Great Expectations, TensorFlow Data Validation, and Pandera provide declarative schemas with automated checks. Distribution monitoring tracks feature distributions over time, catching drift that may not violate a type constraint but still indicates an upstream change whose effect on model performance requires evaluation. Cross-field validation adds the business logic that neither schema nor marginal distributions can express. If `country="USA"`, the `zip_code` should match US postal formats. If `order_status="shipped"`, a `shipping_date` should exist. Age derived from birthdate should match any explicit age field.

12.4.7.1 Debugging data quality incidents

A recommendation model's click-through rate (CTR) drops 8 percent over 5 days. The initial hypothesis focuses on model drift, but the constructed investigation reveals a feature computation defect.

The investigation starts with the model symptom and then works backward through the data path. Model metrics confirm that CTR declined from 4.2 percent to 3.87 percent, but freshness checks show that features still arrive within the 1-hour SLO. That weakens the stale-pipeline hypothesis only if source progress, event-time coverage, and the affected entities are also verified; refresh age alone cannot exclude stale inputs. The investigation therefore tests semantic change as well. Distribution analysis then shows the mean of `user_engagement_score`, the shared feature whose redefinition section 12.4.4 used to motivate versioning, shifting from 0.42 to 0.31, a 26 percent decline. Lineage tracking ties the shift to an upstream pipeline version deployed 5 days earlier.

The shift and deployment history make the new pipeline a candidate cause; a changed population could also move the mean. Replaying the same source records through both pipeline versions with identical event cutoffs isolates the computation. In this scenario, the new version uses a different aggregation window even though its declared feature definition is unchanged, while the previous version matches that definition. This same-record comparison establishes a feature computation defect. The defect is the same class of silent semantic change that feature versioning is meant to name before release; the marginal distribution alone would not establish it.

Assuming revenue falls in proportion to CTR and traffic is uniform over the interval, the modeled impact is 8 percent of $15M/week over 5 days, producing $857K in lost revenue. If automated alerts and effective containment shorten exposure to 4 hours, the modeled impact falls by 97 percent. The resolution restores the previous pipeline and compatible model-feature state. A distribution validation gate then blocks unexplained feature shifts exceeding the scenario's 10 percent threshold; semantic tests remain necessary for bugs that do not cross that band.

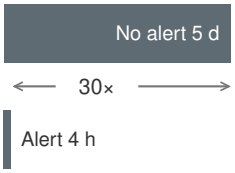

Four-hour alerts shrink a five-day exposure window 30×.

12.4.7.2 Operational integration

Data quality operations integrate into production systems by turning checks into enforceable contracts. **Data contracts** establish formal agreements between data producers and consumers, spec-

ifying schema requirements, freshness SLOs, quality thresholds for completeness and accuracy, and escalation procedures for violations. Release controls enforce those agreements by checking proposed schema or computation changes and requiring consumer review when compatibility cannot be established.

The contract becomes operational through continuous validation and incident response. Ingestion checks reject source data that violates declared constraints, transformation checks test computed values, and serving gates block features known to be invalid. When a check fails, alerts identify severity and affected systems, circuit breakers contain the bad path, and quarantined batches preserve evidence for investigation. A fallback may use cached or last-known-good features only while they remain acceptable under the consuming model's freshness and compatibility contract. Root-cause analysis then repairs the pipeline rather than only clearing the current alert.

At enterprise scale with thousands of features, individual checks need a shared routing and grouping policy so that one upstream failure does not page every downstream consumer separately. Grouping by source system connects symptoms of a shared outage, such as a payment-processing failure that corrupts several features at once, while grouping by computation pipeline directs investigation toward shared transformation code or dependencies. Because streaming and daily batch features operate under different freshness contracts, update-frequency grouping separates their alert thresholds, and business-domain grouping routes the resulting evidence to the teams that understand the affected models. Together, these groupings cut duplicate notifications without removing feature-level checks or suppressing an urgent contract violation.

Freshness monitoring applies the same aggregation rule to time. For 10,000 features updated at different frequencies, checking every feature value can consume substantial capacity. Sampling 1 percent of representative entities reduces that work, but the sample still grows with the entity population and must cover the relevant sources and partitions. Pipeline-level aggregation can reduce checks further when one pipeline updates 500 features and its source-progress watermark covers the group. A recent successful update alone cannot rule out a stalled partition or missing entities. With a bounded number of checks per feature, monitoring can maintain $\mathcal{O}(N_{\text{features}})$ overhead, but that bound trades exhaustive coverage for aggregate or sampled evidence. Critical features need checks whose coverage and detection delay match their failure consequences.

Feature definitions, historical availability, and freshness now belong to the same governed state as model artifacts. A release pipeline must validate that state as a unit, retain a compatible rollback target, and limit live exposure while evidence accumulates.

12.5 CI/CD for ML at Scale

Versioned model dependencies and feature contracts define a candidate release, but they do not establish that it is safe to expose. Each model update must navigate the dependency web. In the e-commerce graph of figure 12.4, an embedding model update might require re-evaluation of four transitive downstream components, including the business rules layer, before any component can safely reach production. This coordination challenge transforms continuous integration and continuous deployment from a per-model concern into a platform orchestration problem. Whereas traditional software CI/CD focuses on code correctness and deployment reliability across interacting components, ML CI/CD adds learned behavior and changing data to that contract, requiring automated pipelines that validate data, verify statistical model performance, and confirm that upstream changes preserve downstream quality across every intermediate deployment state.

12.5.1 Training pipeline automation

When a training run is kicked off from an engineer's laptop with unrecorded data paths and ad-hoc hyperparameters, no reviewer can reproduce the resulting checkpoint or verify what changed when it fails in staging. CI/CD for machine learning begins by making each training run an accountable producer of deployable artifacts. The pipeline records versioned inputs and execution conditions, defines recovery behavior, and preserves the results needed for deployment validation. Automation makes those obligations repeatable; reproducibility still depends on the recorded data, code, environment, and sources of nondeterminism.

12.5.1.1 Pipeline stages

Figure 12.7 stages the delivery lifecycle across continuous integration, training, and deployment, with an evaluation gate that blocks or registers a candidate. Monitoring returns drift signals to data validation. A confirmed population change may justify retraining, while a corrupted feature pipeline requires repair before it supplies another training run. The sequence matters because each stage establishes evidence for the next rather than treating every alert as permission to train and deploy.

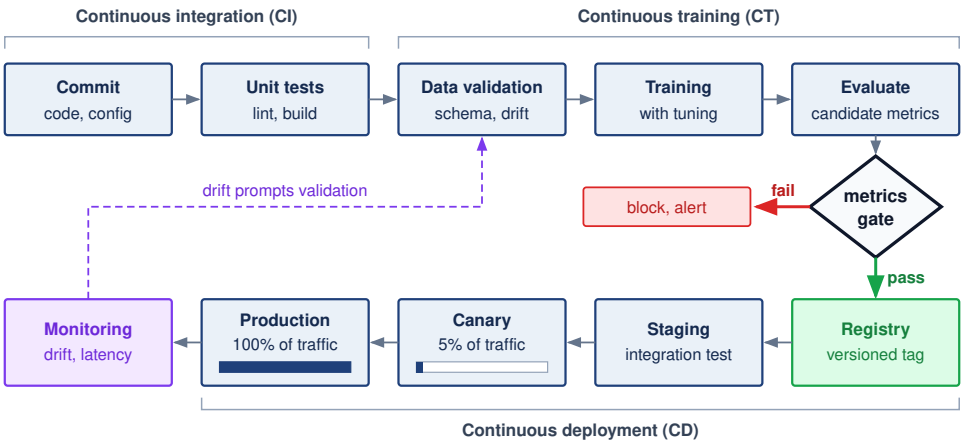

Figure 12.7: ML CI/CD Pipeline: Read the top row left to right and the bottom row right to left. Continuous integration checks each commit; continuous training validates data, trains, and evaluates against a metrics gate that either blocks the candidate and raises an alert or registers it. Continuous deployment then moves the registered version through staging, a canary slice of traffic, and full production. Monitoring closes the loop: detected drift returns to data validation and investigation before a decision to retrain.

1. Data validation pins the data version and rejects schema, freshness, or distribution failures before they consume training capacity.
2. Feature engineering preserves the feature contract between training and serving, so preprocessing changes do not become silent production regressions.
3. Training records code, data, hyperparameters, hardware envelope, and random seeds so a promising artifact can be reproduced or challenged later.
4. Evaluation compares the candidate with baselines, slice metrics (quality measured separately on subsets such as a region or user cohort), latency budgets, and task-specific gates before any user traffic is exposed.
5. Artifact generation packages the model with its serving configuration, dependency versions, and rollback target.
6. Registration records the artifact in the model registry with lineage, approvals, and deployment eligibility.

Each stage needs a defined retry boundary and durable outputs. If evaluation fails, the pipeline can reuse completed data-validation and feature-engineering results when their data, code, configuration, and relevant execution environment remain unchanged. Missing outputs or a changed dependency invalidate that reuse. Idempotent publication prevents a retry from registering duplicate artifacts or repeating external side effects.

12.5.1.2 Pipeline orchestration

Training pipelines require orchestration systems that make dependency order, resource allocation, and failure recovery explicit. The orchestrator represents the workflow as a directed acyclic graph (DAG) and tracks dependencies between stages. It retries transient failures without rerunning

unaffected work, schedules scarce resources such as GPUs and memory, caches intermediate results when inputs are unchanged, and records logs and artifacts for later debugging.

Common orchestration choices include Kubeflow Pipelines[6] (Bisong 2019), Airflow with ML extensions, and cloud-native solutions like Vertex AI Pipelines or SageMaker Pipelines. The choice depends on which dependency and resource contracts the platform must enforce, as well as existing infrastructure, team expertise, and scale requirements. Orchestration fixes the execution contract; parameterization fixes the variation contract, so the same graph can run against different data, models, and hardware envelopes without becoming a new program each time.

[6] **Kubeflow**: Kubeflow Pipelines connects workflow dependencies to Kubernetes execution and resource requests. Declared requests let scheduling account for resources needed by a step, but do not eliminate contention on shared storage, networks, or other services. Platform policy must therefore combine admission and resource allocation with measurements of actual execution behavior.

12.5.1.3 Pipeline parameterization

Effective pipelines separate configuration from code (listing 12.3). The listing's evaluation block names two standard ranking metrics: normalized discounted cumulative gain at rank 10 (NDCG@10), which rewards placing relevant items near the top of a list, and mean reciprocal rank (MRR), which averages the inverse position of the first relevant item.

Listing 12.3: Pipeline Parameterization: A YAML configuration separating data paths, feature sets, training hyperparameters, and evaluation criteria from pipeline code.

```yaml
training_pipeline:
  model_type: transformer_ranking
  data:
    train_path: gs://data/train/2024-01-*
    eval_path: gs://data/eval/2024-01-15
    schema_version: v3.2
  features:
    user_features: [embedding, history, demographics]
    item_features: [embedding, attributes, popularity]
  training:
    epochs: 10
    batch_size: 4096
    learning_rate: 0.001
    optimizer: adam
    hardware: 4xA100
  evaluation:
    metrics: [ndcg_10, mrr, coverage]
    baseline_model: ranking_v2.1.0
```

The separation enables configuration-driven training because the platform can vary data versions, hyperparameters, and resource requests without rewriting pipeline code. A parameter file is also an audit input, but a path or version label alone does not prove which bytes a run consumed. The execution record must resolve those references to immutable data and artifacts and retain the actual environment, feature definitions, and evaluation policy. That record lets a release reviewer reconstruct the conditions behind the candidate's evidence.

12.5.2 Validation gates

An immutable execution record tells a reviewer what a training run produced. Validation gates decide whether that candidate may advance toward production, because promoting a model solely for a 1 percent aggregate NDCG gain can still double p99 tail latency, degrade a protected subgroup, or rely on a corrupted upstream feature. A release gate consumes a versioned evidence bundle tying the candidate and its dependencies to metrics, thresholds, and accountable owners. Model performance, latency, policy and fairness constraints, data quality, and operational readiness each answer a different failure mode. Together they form an evaluation contract, a single versioned record of release evidence that remains attached to the release as it moves from offline validation to production experiments (section 12.5.8).

12.5.2.1 Performance gates

Performance validation combines minimum acceptable performance with comparisons against the production baseline and historical results. The release policy names the primary improvement

sought and the regression bounds permitted on secondary metrics, rather than requiring every metric to improve simultaneously. A performance gate therefore checks three conditions: absolute thresholds, relative improvement over the production model, and regression bounds on secondary metrics (listing 12.4).

Listing 12.4: Performance Gate Evaluation: A validation function that checks absolute thresholds, relative improvement over the production model, and regression bounds on secondary metrics.

```
def evaluate_performance_gate(
    candidate_metrics, production_metrics, thresholds
):
    """
    Evaluate whether candidate model passes performance gates.

    Returns tuple of (passed: bool, reasons: list)
    """
    import math

    required_metrics = ("ndcg_10", "mrr", "coverage")
    required_thresholds = ("min_ndcg", "min_improvement", "max_regression")
    try:
        values = [candidate_metrics[k] for k in required_metrics]
        values += [production_metrics[k] for k in required_metrics]
        values += [thresholds[k] for k in required_thresholds]
        valid = all(math.isfinite(v) for v in values)
        valid = valid and production_metrics["ndcg_10"] > 0
        valid = valid and 0 <= thresholds["max_regression"] <= 1
    except (KeyError, TypeError, ValueError):
        valid = False
    if not valid:
        return False, ["Missing, nonfinite, or invalid gate inputs"]

    reasons = []

    # Absolute threshold check
    if candidate_metrics["ndcg_10"] < thresholds["min_ndcg"]:
        reasons.append(
            f"NDCG_10 {candidate_metrics['ndcg_10']:.4f} below minimum {thresholds['min_ndcg']}"
        )

    # Relative improvement check
    relative_improvement = (
        candidate_metrics["ndcg_10"] - production_metrics["ndcg_10"]
    ) / production_metrics["ndcg_10"]
    if relative_improvement < thresholds["min_improvement"]:
        reasons.append(
            f"Improvement {relative_improvement:.2%} below minimum {thresholds['min_improvement']:.2%}"
        )

    # Regression check on secondary metrics
    for metric in ["mrr", "coverage"]:
        if candidate_metrics[metric] < production_metrics[metric] * (
            1 - thresholds["max_regression"]
        ):
            reasons.append(
                f"{metric} regression exceeds {thresholds['max_regression']:.2%} tolerance"
            )

    return (len(reasons) == 0, reasons)
```

12.5.2.2 Latency gates

Production models must meet workload-specific latency requirements. Validation measures representative request shapes at expected load, batching, and hardware. Illustrative thresholds span two orders of magnitude, from a 5 ms median for fraud detection to a 500 ms median time to

first token (TTFT) for an LLM, and they are policy choices rather than universal model-type requirements (table 12.12). The LLM row measures TTFT, reflecting the prefill phase of section 10.8; completion latency requires a separate output-length contract for decode. The other rows measure request-to-response latency.

Table 12.12: Illustrative Latency Gates: Assumed p50 and p99 service targets under a declared request distribution and load. LLM targets refer to TTFT; other targets cover the complete response. Product risk and the service contract determine thresholds and exception authority.

Model type	p50 target	p99 target	Gate action if exceeded
LLM (TTFT)	500 ms	2000 ms	Block deployment, require optimization
Recommendation	10 ms	50 ms	Block deployment
Fraud Detection	5 ms	20 ms	Block deployment, high priority
Vision	50 ms	200 ms	Warning, conditional approval

Latency gates catch overt performance violations, but some regressions pass every gate and still degrade the product. The most dangerous failures are semantically silent: valid outputs built on corrupted inputs.

Example 12.2: Silent model regression

Scenario: A high-throughput product ranking model (25,000 QPS) passes offline validation and serving latency gates (per-request latency $L_{\text{lat,SLO}} < 12$ ms), but causes a 3.2 percent conversion rate drop post-deployment.

Diagnosis: An upstream schema change caused a key user-affinity feature to return `null` for 8 percent of queries. The serving layer silently imputed `null` as `0.0`, a valid float value, so HTTP error rates and latency remained normal while prediction quality degraded unalerted for 48 hours.

Systems lesson: Syntactically valid feature imputations cause semantically silent production failures. Automated release pipelines and monitoring systems must evaluate feature distribution shifts and null rates alongside system uptime and latency SLOs.

12.5.2.3 Fairness gates

Fairness concepts become operational at release time when policy choices are encoded as thresholds. For models affecting users, fairness validation turns a policy commitment into a release gate. A gate cannot test whether "the model is fair" in the abstract, because different contexts encode fairness differently and some definitions cannot hold at once. Each chosen definition becomes an executable bound on a measured gap between protected groups. A demographic-parity gate bounds the difference between groups in the rate of positive predictions; an equalized-odds gate bounds the difference in prediction rates between groups that share the same true outcome (Hardt et al. 2016). The operational point is not that every gate should enforce every definition. The gate must record which definition the organization chose, compare the measured value against historical baselines as well as absolute thresholds, and route borderline cases to human review because a near-threshold fairness result is also a governance decision.

↪ The fairness criteria behind these gates, and why they cannot all hold at once, are developed in section 16.5.1.

12.5.2.4 Data quality gates

Before training or deployment, data quality validation reapplies the schema, distribution, freshness, and completeness checks of section 12.4.7 as a release gate (Caveness et al. 2020). Passing them establishes evidence for the properties tested; it does not rule out an untested semantic error.

These gates complete the offline evidence bundle. The next evidence must come from production traffic, gathered first without exposing users and then at a bounded share of live requests.

12.5.3 Production experimentation as a release gate

Offline validation on static holdouts cannot catch a candidate model that doubles tail latency under live concurrency or reduces user clicks when its predictions interact with real traffic. Once the offline gates pass, production experiments evaluate the candidate under a controlled exposure policy, starting from shadow traffic, which exposes no user to its outputs. Shadow traffic tests behavior without returning candidate outputs, interleaving compares rankers within a combined presentation, and A/B tests compare assigned variants. Network-effect checks determine whether interactions between experimental units undermine that comparison. These methods answer different questions; every release need not use all of them. Each supplies evidence for the online-experiment plane of the evaluation contract assembled in section 12.5.8, so the statistics that follow are tools for filling one row of the release record, not a separate discipline.

Before exposure begins, the release record must name the invariant metrics, decision threshold, maximum user and infrastructure blast radius, stopping rule, rollback owner, and evidence required for promotion. A candidate that improves its primary metric but exhausts latency, reliability, fairness, or cost headroom fails the gate.

12.5.3.1 Shadow deployment and traffic replay

Shadow deployment runs the new model alongside production, receiving copied inputs and logging outputs without returning them to users. Figure 12.8 contrasts mirroring to an isolated candidate (Panel A) with canary deployment (Panel B), which exposes a controlled share of traffic to candidate responses. Shadow execution can reveal defects absent from offline tests, but its resource consumption and external side effects still require containment. It also cannot reveal how users would react to outputs they never receive.

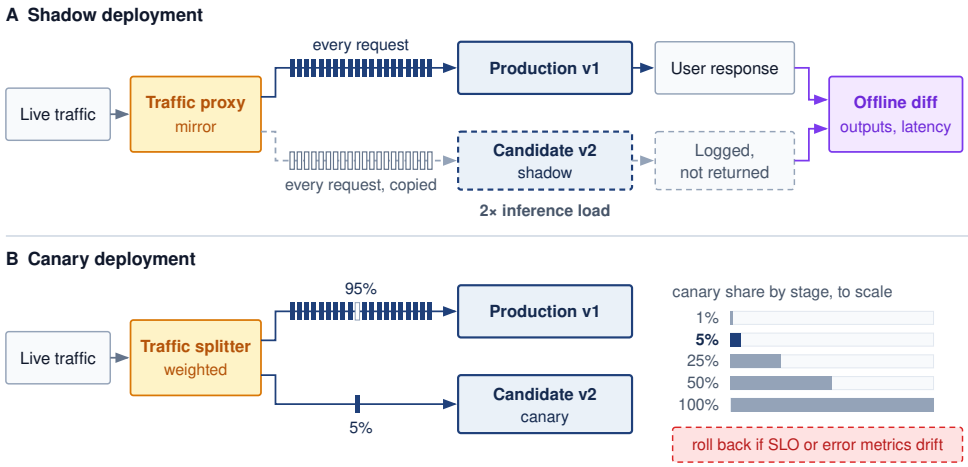

Figure 12.8: Shadow and Canary Deployment: In shadow deployment (A), the proxy mirrors all production traffic to the candidate; users receive only the production response, and both outputs are compared offline. No user sees the candidate's output, but mirroring every request doubles inference load, so the shadow path needs its own capacity and must be isolated so it cannot slow production. In canary deployment (B), the splitter sends a small weighted slice of live traffic to the candidate (here the 5 percent stage) and widens it through the 1, 5, 25, 50, and 100 percent stages as evidence accumulates, rolling back if SLO or error metrics drift.

Shadow deployment earns its duplicate infrastructure cost when it answers validation questions offline tests cannot. A full-load shadow run measures errors, memory growth, and latency against the release SLOs, exposing leaks and resource contention that small offline datasets miss. Comparison under real traffic reveals output shifts such as changed confidence scores in a classifier or changed diversity in a recommender. That evidence holds only for the observed workload, duration, dependencies, and contention, and rare inputs may never arrive, so the shadow load profile and run

length belong in the release record beside the result. Output differences require interpretation; they do not by themselves establish improvement or regression.

Shadow deployment requires capturing and replaying production traffic, and the replay choice determines the trade-off among fidelity, cost, and freshness. Three architectural patterns address different operational requirements.

Live mirroring duplicates every production request to the shadow model in real time. The production model serves the response while the shadow model processes the same input in parallel, providing immediate validation at full production scale but requiring shadow infrastructure capable of handling 100 percent traffic load. To keep this validation path from becoming a production dependency, the serving system must invoke shadow requests asynchronously, enforce timeouts, shed shadow load when it falls behind, and isolate resources so shadow work cannot affect production responses.

Sampled replay mirrors a configurable percentage of production traffic to shadow models. This reduces infrastructure costs, but the sampling rate and observation window must still supply enough eligible examples for the targeted failures. Uniform random sampling estimates the sampled traffic population, while stratification can ensure coverage of important segments. Adaptive sampling concentrates investigation on patterns where candidate and production outputs diverge. Such targeted samples need recorded inclusion probabilities, the chance that each request entered the sample, and an estimator that corrects for them before their error rates can be interpreted as fleet-wide rates.

Batch replay captures production traffic logs and replays them asynchronously against shadow models. This decouples validation from the live request path and supports regression tests during available capacity windows. Fidelity depends on what the logs preserve. Replaying request contents without their historical feature state can change the task, while replaying them faster changes queueing and load. A load test must reproduce the relevant arrival pattern and serving conditions; historical replay still cannot observe user responses to candidate outputs.

Effective shadow deployment requires quantitative comparisons tied to the intended change. Output divergence locates differences from production: classification systems track disagreement rates, probability shifts, and class-specific concentration, while regression systems can compare prediction errors or differences between outputs. A changed prediction is not inherently a regression; outcome evidence or a validation oracle, an independent check that determines the correct output, must establish its quality. Performance metrics assess latency, throughput, and resource consumption against the operating contract. A candidate with equivalent accuracy but 50 percent higher p99 latency needs review against that contract, even if both versions remain below the SLO.

The metric set must also separate operational failures from genuine model changes. Error-mode metrics count timeouts, exceptions, malformed outputs, and null predictions. A shadow model that times out on 0.1 percent of requests encounters 1,000 failures per day at 1M requests/day scale, and the request patterns that trigger those failures guide remediation. Statistical validation then determines whether observed differences represent genuine model changes or random variation. Because production and shadow models process the same requests, their outcomes are paired. Define a binary failure against a labeled outcome or validation oracle, let n_{prod} count requests on which production fails and shadow succeeds, and let n_{shadow} count requests on which shadow fails and production succeeds. McNemar's large-sample test uses the discordant pairs:

$$\chi^2 = \frac{(n_{\text{prod}} - n_{\text{shadow}})^2}{n_{\text{prod}} + n_{\text{shadow}}}$$

The test starts from a null hypothesis, the assumption that production and shadow fail with equal probability. Under that hypothesis, the statistic approximately follows a chi-square distribution with one degree of freedom. The significance level α_{sig} is the probability of rejecting the null hypothesis when it is true, and the critical value is the threshold that the statistic exceeds with that probability when the null hypothesis holds.

Equivalently, the p-value is the probability of a statistic at least as large as the one observed if the null hypothesis holds, and the test rejects when the p-value falls below α_{sig}. For a one-degree-of-freedom test at $\alpha_{\text{sig}} = 0.05$, the critical value is 3.84, so reject equal failure probabilities when $\chi^2 > 3.84$. A rejection indicates that the failure probabilities differ, and the larger discordant count

shows which model fails more often. Marginal failure rates, such as 1 percent for production and 1.5 percent for shadow, do not determine this statistic by themselves; the paired discordant counts n_{prod} and n_{shadow} must be retained.

12.5.3.2 Shadow deployment workflow

Consider an illustrative fraud detection service processing 5 million transactions daily. A sampled shadow phase mirrors 10 percent of traffic for three days. Assume the eligible labeled subset yields recall of 94.2 percent for the shadow model versus 94.5 percent in production, and precision of 87.1 percent versus 82.3 percent. These point estimates alone establish neither improvement nor noninferiority, the finding that the shadow model is worse by no more than an agreed tolerance. The gate also needs labeled counts, paired outcomes, label maturity (enough elapsed time for each transaction's true outcome to be known), uncertainty intervals, and a prespecified recall tolerance. Assume shadow p99 latency is 45 ms versus 38 ms in production, both below this service's 50 ms SLO, a looser contract than the illustrative 20 ms fraud-detection gate in table 12.12. The team proceeds to full shadow only after the statistical gate passes.

The full-shadow phase mirrors 100 percent of traffic for five days. In this scenario, the same labeled-data analysis supports the precision gain at full scale, but the run also exposes an edge case. The shadow model flags 0.02 percent of transactions as errors due to an unexpected feature distribution, which is 1,000 transactions per day. The root cause is heightened sensitivity to outliers in transaction amount. Adjusting feature clipping thresholds reduces the error rate to 0.001 percent in a new shadow run. Because clipping also changes model inputs, the revised candidate must pass the precision, recall, latency, and slice gates again before canary deployment.

Post-deployment validation then compares actual production metrics with the shadow predictions. Precision materializes at 87.3 percent in production vs. 87.1 percent in shadow, a nearby point estimate. Confirming agreement still requires uncertainty bounds and checks for label or traffic shifts; similarity between two percentages is insufficient.

12.5.3.3 Criteria for shadow deployment

Shadow deployment earns its cost when the additional evidence reduces expected deployment harm enough to justify duplicate execution, comparison logging, and analysis. That cost includes infrastructure that bounds the shadow path's effect on production. The mirroring layer enforces sampling, timeouts, and load shedding. Separate compute pools, network allocation, and database capacity keep shadow work from stealing production capacity, because shared queues or dependencies can still transmit failures that asynchronous invocation alone does not prevent.

New architectures, consequential predictions, live feature dependencies, and tight latency budgets can make the evidence worth that cost, but only for failures shadow traffic can reveal; user responses to candidate outputs remain unobserved, and rare events may remain absent. A smaller shadow phase, or a direct move from offline validation to a bounded canary, can be appropriate when existing evidence covers the relevant failure modes and residual exposure is acceptable. Routine retraining does not establish low risk by itself, because changed data or parameters can alter behavior without an architecture change. The platform should compare the additional evidence shadow execution can provide with its resource cost, the harm a canary could cause before containment, and the time needed to observe that harm.

12.5.3.4 Interleaving experiments

Interleaving experiments[7] were developed for ranking and search evaluation (Chapelle et al. 2012) and are also used in recommender personalization experiments (Netflix Technology Blog 2017). Rather than splitting users between variants, interleaving presents items from both variants to each user, then measures which items users engage with.

The advantage of interleaving is statistical efficiency. An interleaving experiment can require far fewer samples than A/B testing, and Netflix reports more than 100× fewer users for some ranking experiments (Netflix Technology Blog 2017), because each user provides direct comparison signals rather than contributing only to aggregate statistics.

In an interleaving workflow, the variants produce ranked lists that a method such as team draft or probabilistic interleaving combines into one presentation. The mixing and credit-assignment

[7] **Interleaving experiments**: Originally developed for search engine evaluation, interleaving merges results from two rankers into one list shown to each user, then credits clicks to the originating ranker (Chapelle et al. 2012). For fleet-scale recommendation systems iterating on many model variants, the sample efficiency of that within-user comparison translates directly into faster experiment cycles and lower opportunity cost from suboptimal models in production.

rules must control position effects and identify which ranker contributed each item. Engagement with those items supplies a relative preference signal. Figure 12.9 contrasts the between-cohort comparison of A/B testing (Panel A) with this within-user comparison (Panel B). The latter can screen candidates efficiently, but preference within a blended list does not directly measure the effect of deploying either complete ranking policy on long-term outcomes such as retention. Those outcomes still require an experiment designed to measure them.

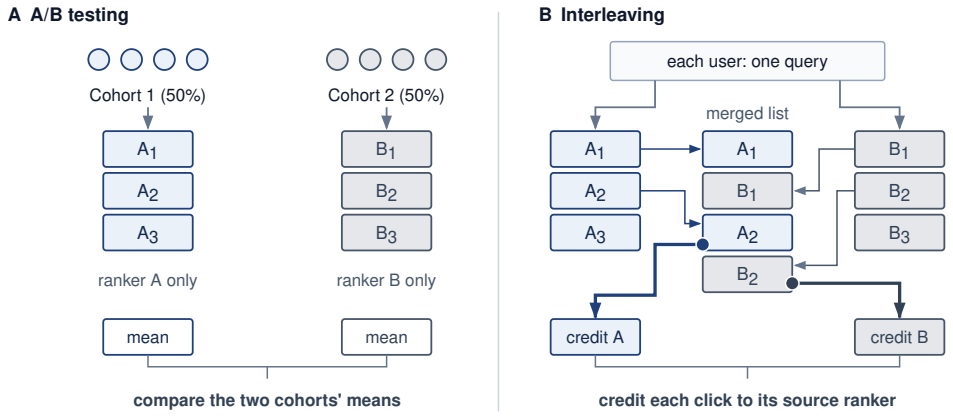

Figure 12.9: Interleaving vs. A/B Testing: (A) A/B testing splits users into cohorts, and each cohort sees only one ranker's list, so the comparison is between cohort means and must overcome the variation between users. (B) Interleaving has both rankers rank the same query for each user and merges the results into one list (here by team draft); each click is credited to the ranker that contributed the item, so every user supplies a paired, within-user comparison. That pairing can make ranker-preference screening more sample-efficient; it does not measure the full deployment outcome assessed by an A/B test (Netflix Technology Blog 2017).

12.5.3.5 A/B test sizing and significance

At platform scale, experiment design determines how long a release occupies capacity and exposes users. A/B testing provides frameworks for comparing model variants (Kohavi et al. 2009, 2020), but its release decision depends on statistical power,[8] significance thresholds, and multiple testing correction. The formulas below supply inputs to the operational contract: the number of eligible observations, the promotion rule, and the permitted stopping schedule. They do not establish that behavior remains reliable outside the evaluated population.

For an equally allocated comparison of two proportions, the standard normal-approximation sample plan uses four parameters: significance level (α_{sig}), statistical power $(1-\beta_{\text{stat}})$, baseline conversion rate (p_0), and minimum detectable effect (δ_{MDE}). Equation 12.4 estimates observations per variant, assuming independent eligible outcomes and a small absolute difference near the baseline rate. Holding the other inputs fixed, the requirement scales inversely with the square of that difference:

$$n = \frac{(z_{\alpha/2} + z_{\beta})^2 \times 2p_0(1-p_0)}{\delta_{\text{MDE}}^2} \tag{12.4}$$

where $z_{\alpha/2}$ is the two-sided critical value for significance level α_{sig} (typically 1.96 for $\alpha_{\text{sig}} = 0.05$), z_{β} is the critical value for power (typically 0.84 for 80 percent power), p_0 is the baseline rate, and δ_{MDE} is the minimum detectable effect as an absolute difference. Both critical values use the positive upper-tail convention: the upper-tail areas are $\alpha_{\text{sig}}/2$ and β_{stat}, respectively.

Once data is collected, a two-proportion z-test determines if the observed difference is statistically significant. Equation 12.5 computes the test statistic as the difference in observed rates normalized by the pooled standard error:

[8] **A/B testing power**: The probability that an experiment correctly detects a true effect ($1-\beta_{\text{stat}}$, typically set to 0.8). At fleet scale, the systems cost of power is exposure time. If the expected improvement is small (for example, 0.1 percent), the experiment may require millions of samples to reach statistical significance, consuming substantial serving resources (C_{infer}) for weeks.

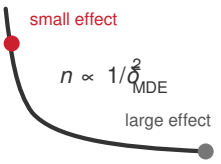

Detecting smaller effects explodes the required sample size.

$$z = \frac{\hat{p}_B - \hat{p}_A}{\sqrt{\hat{p}(1-\hat{p})(\frac{1}{m_A} + \frac{1}{m_B})}} \tag{12.5}$$

where $\hat{p}_A$ and $\hat{p}_B$ are the observed conversion rates for control and treatment, m_A and m_B are sample sizes, and $\hat{p} = \frac{m_A\hat{p}_A + m_B\hat{p}_B}{m_A + m_B}$ is the pooled proportion. For the two-sided test, if $|z| > z_{\alpha/2}$, reject the null hypothesis and conclude the variants differ significantly.

12.5.3.6 Sample size for recommendation models

A recommendation system has baseline CTR of 5 percent. The team wants to detect a 10 percent relative improvement (0.5 percentage points absolute) with 95 percent confidence and 80 percent power.

Parameters:

- $z_{\alpha/2} = 1.96$ (95 percent confidence, two-tailed)
- $z_\beta = 0.84$ (80 percent power)
- $p_0 = 0.05$ (baseline CTR)
- $\delta_{\text{MDE}} = 0.005$ (0.5 percentage point improvement)

Calculation:

$$n = \frac{(1.96 + 0.84)^2 \times 2 \times 0.05 \times 0.95}{0.005^2}$$

$$n = \frac{7.84 \times 0.095}{0.000025} = \frac{0.7448}{0.000025} = 29{,}792$$

Each variant requires approximately 30,000 samples, totaling 60,000 observations. If requests supply independent eligible observations with immediately available outcomes, 1 million requests per day split evenly between variants collects that sample in less than 2 hours. However, for a model with 1 percent baseline CTR detecting a 5 percent relative improvement (0.05 percentage points), the calculation yields:

$$n = \frac{7.84 \times 2 \times 0.01 \times 0.99}{0.0005^2} = \frac{0.155232}{0.00000025} = 620{,}928$$

Now each variant needs approximately 621K samples. Under the same independence and outcome-availability assumptions, 1M total requests/day split evenly between the two variants collects that sample in about 30 hours. Required observation windows can extend either experiment beyond its collection time. Here, the lower baseline and smaller relative improvement also imply a much smaller absolute effect. At a fixed relative improvement, rarer outcomes require more observations; a fixed absolute effect would give a different comparison.

12.5.3.7 Multiple testing correction

Running multiple A/B tests simultaneously or sequentially without correction inflates the familywise error rate, the probability that at least one test in the family reports a false positive. With 20 independent tests whose null hypotheses are all true, at $\alpha_{\text{sig}} = 0.05$, that probability is:

$$\Pr(\text{at least one false positive}) = 1 - (1 - \alpha_{\text{sig}})^k = 1 - 0.95^{20} = 0.642$$

The result is a 64 percent chance of at least one false rejection; a two-sided rejection can point in either direction. Two families of correction address this problem.

Bonferroni correction adjusts the significance threshold to $\alpha'_{\text{sig}} = \frac{\alpha_{\text{sig}}}{k}$ for k tests. This is conservative but simple. For 20 tests with $\alpha_{\text{sig}} = 0.05$, use $\alpha'_{\text{sig}} = 0.0025$ for each test. This controls the familywise error rate but reduces statistical power; the Šidák variant, which sets the per-test level so that independent tests reach the familywise target exactly, is only slightly more lenient (0.00256 per test for 20 tests).

False discovery rate (FDR) control using the Benjamini-Hochberg procedure bounds the expected proportion of false discoveries among rejected hypotheses under independence or suitable positive

dependence. It is not a hard cap on the probability of making at least one false discovery in a batch, which makes it appropriate when screening many candidate changes where a small fraction of false positives is acceptable. Order the k observed p-values from smallest to largest, $p_{\text{val}_{(1)}} \leq p_{\text{val}_{(2)}} \leq \ldots \leq p_{\text{val}_{(k)}}$, and find the largest rank r whose p-value satisfies:

$$p_{\text{val}_{(i)}} \leq \frac{r}{k} \times \alpha_{\text{sig}}$$

Reject all hypotheses $H_{(1)}, \ldots, H_{(r)}$ up to that cutoff rank. Only the most extreme p-value $p_{\text{val}_{(1)}}$ faces the strict Bonferroni threshold $\frac{1}{k}\alpha_{\text{sig}}$; each subsequent rank r is compared against a linearly growing allowance $\frac{r}{k}\alpha_{\text{sig}}$. This step-up procedure therefore yields substantially higher statistical power than Bonferroni when evaluating large experiment portfolios.

12.5.3.8 Multiple testing scenario

A platform team runs 30 A/B tests per quarter comparing candidate models. Suppose every null hypothesis is true and each test has a calibrated false-positive probability of $\alpha_{\text{sig}} = 0.05$. Without correction, the expected count is $30 \times 0.05 = 1.5$ false rejections per quarter, or approximately 6 models over a year. These are false claims of a difference, not necessarily improvements: a two-sided rejection can point toward either benefit or harm.

Applying Bonferroni correction: $\alpha'_{\text{sig}} = \frac{0.05}{30} = 0.00167$ per test. This requires larger sample sizes. For the recommendation model example in section 12.5.3.6 (5 percent baseline, 0.5-percentage-point effect), the original requirement was 30K samples per variant. With a two-sided Bonferroni threshold, the critical value rises to about 3.14, increasing the requirement to roughly 60K samples per variant.

Using FDR control at $q = 0.05$, the Benjamini-Hochberg procedure bounds the expected false-discovery proportion across repeated experiment families under its dependence assumptions, counting the proportion as zero when there are no rejections. Observing 10 significant results out of 30 tests does not establish how many of those 10 are false or give a conditional guarantee for that quarter. The procedure accepts a different error criterion from familywise control and can gain power in exchange.

The choice begins with the error the decision must control. When even one false promotion among the planned comparisons has serious consequences, familywise control is appropriate; Bonferroni is one simple option. Exploratory screening may instead use FDR control when a bounded expected proportion of false discoveries is acceptable and later validation will assess candidates individually. The experiment family and error criterion belong in the release plan before results are inspected.

> Checkpoint 12.3: Family-wise error rate under multiple testing
>
> The correction methods are Bonferroni and Benjamini-Hochberg. Apply the family-wise reasoning to a smaller platform.
>
> - ☐ **Family-wise error**: for 20 simultaneous A/B tests per quarter at $\alpha_{\text{sig}} = 0.05$ with no correction, calculate the expected number of false positives and the family-wise error rate under independence.
> - ☐ **Bonferroni correction**: apply Bonferroni correction, calculate the per-test threshold α'_{sig}, compute the family-wise error rate after correction, and compare it to the uncorrected rate.
> - ☐ **Bonferroni vs. FDR**: for 20 tests where the team expects most candidates to genuinely improve the metric, choose Bonferroni or Benjamini-Hochberg FDR control and justify the statistical-power trade-off.

12.5.3.9 Sequential testing and early stopping

Traditional A/B tests fix sample size in advance and evaluate once. Release teams rarely wait that long without looking, and looking is where the error rate leaks. A fixed-sample test checked after every batch of observations and stopped at the first $|z| > 1.96$ gets many chances to cross the threshold by chance, so its realized false-positive rate climbs well above the nominal 5 percent.

Sequential testing budgets for those looks in advance, allowing monitoring during data collection with principled early stopping rules. Stopping can occur earlier when the evidence is strong, but the gain depends on the hypotheses, effect size, and stopping design.

The sequential probability ratio test, the classic design, updates a likelihood ratio after every observation and stops when the ratio crosses either of two thresholds set from the target error rates. Because it assumes two fully specified hypotheses, release experiments more often use group sequential designs. Group sequential methods divide the experiment into planned analysis stages. At each stage, the test statistic is compared with adjusted thresholds that maintain the overall α_{sig}. The boundary family decides how that error budget is spent across the looks. An O'Brien-Fleming boundary spends almost none of it early, demanding overwhelming evidence to stop at the first looks and keeping the final threshold close to the fixed-sample value. A Pocock boundary spends it evenly, applying the same stricter threshold at every look. For a canary whose guardrail metric is checked every hour, the evaluation contract must prespecify a valid sequential stopping rule, such as one of these boundaries. Comparing each hourly look against the single-look 1.96 reintroduces the inflation the design exists to prevent.

A valid stopping rule still rests on the independence and stationarity assumptions behind these formulas, and production tests can violate both. Carryover occurs when treatment changes behavior after exposure ends, so a design that reassigns users must account for that persistence. Novelty can make early responses differ from later behavior, and no fixed number of weeks guarantees that the effect has settled, so observation windows must cover the relevant behavioral cycles and outcome delays. Surrogate metrics such as clicks can also move in a different direction from long-term objectives such as retention and revenue, which is why a release gate pairs the short-term surrogate with guardrail metrics that may require longer observation. Interference between users is the most consequential violation for social and marketplace products, because it biases the estimated effect itself rather than only its variance.

12.5.4 SUTVA violations and network effects

When User A on a social platform receives better content recommendations from a treatment model and shares those posts with friends in the control group, User B's engagement rises even though User B never received the new algorithm. The control condition becomes contaminated through the natural mechanics of a networked product, violating the assumption behind every sample-size and significance calculation above that one user's treatment assignment does not change another user's outcome.

The stable unit treatment value assumption (SUTVA) requires well-defined treatment versions and no interference between units (Rubin 1980). Its no-interference component says that user i's outcome $Y_i^{\text{pot}}(w)$ depends on their own assignment w, not on other users' assignments. Versioned releases address the treatment-definition requirement; network interactions challenge the second requirement. Ignoring those interactions can bias the effect estimate used for a deployment decision (Hudgens and Halloran 2008; Eckles et al. 2017).

Interference can travel through direct connections, shared markets, or geographic and temporal spillovers. A treated user may share recommended content with control users, changing their outcomes even though they never receive the new ranker. A ridesharing incentive may change driver supply and wait times for both groups. Local merchant recommendations may redirect traffic from neighboring areas assigned to control. In these examples, independent user assignment remains possible, but each user's assignment no longer describes all the treatment exposure affecting that user's outcome. The direction and size of the resulting bias depend on the interaction and the effect being estimated.

12.5.4.1 Quantifying SUTVA violations

Grouping connected users into randomized clusters changes the precision calculation as well as the exposure pattern. For clusters of equal or similar size with a common within-cluster outcome correlation, equation 12.6 gives a planning approximation for variance inflation relative to independent observations:

$$\text{DEFF} \approx 1 + (m-1)\rho_{\text{ICC}} \tag{12.6}$$

where m is the average cluster size and ρ_{ICC} is the intra-cluster correlation of outcomes (how similar outcomes are within connected user groups). Multiplying an independent-observation sample plan by this factor approximates the additional observations needed under those assumptions. Unequal cluster sizes and the chosen estimator require a more specific precision analysis. This variance adjustment does not identify or remove bias from interference across clusters.

12.5.4.2 Network effect bias in social recommendation

An illustrative social platform tests a new feed ranking algorithm. Individual user randomization assigns 50 percent of users to treatment.

Experimental setup:

- 10 million users, average 150 connections each
- Outcome: daily engagement minutes

With independent 50 percent assignment, a typical control user's 150 connections split roughly evenly between arms, so almost every control user is exposed to treatment through dozens of treated neighbors.

Naive analysis results:

- Treatment group: 45.2 minutes average
- Control group: 43.8 minutes average
- Measured effect: +1.4 minutes (+3.2 percent)

Network analysis confirms that control users have on average about half of their connections in treatment. Those treated connections share algorithmically boosted content that control users see, inflating control group engagement. The observed data cannot say by how much, because almost no control user is unexposed. The sensitivity analysis therefore supplies the missing quantity as an explicit assumption, an unexposed control mean of 42.3 minutes:

- Adjusted control baseline: 42.3 minutes (what control would show without spillover)
- Contrast against the assumed baseline: +2.9 minutes (+6.9 percent)
- Under this assumption, spillover inflated control by 1.5 minutes, halving the measured effect

The assumed counterfactual is not identified by the observed group means or the connection count. Estimating an exposure-specific effect requires an explicit exposure mapping, assignment probabilities, sufficient probability of each exposure, and an estimator matched to those choices. The arithmetic shows how a plausible spillover assumption can change the interpretation of an experiment. It does not recover a known true effect, and a design effect for cluster randomization is a separate variance calculation that cannot correct interference bias.

12.5.4.3 Detection strategies

Network-exposure diagnostics test for evidence of interference. Ego-network analysis summarizes each control user's treated neighbors as $e_i = |\{j \in \mathcal{N}_i : w_j = 1\}|/|\mathcal{N}_i|$, where $\mathcal{N}_i$ is the recorded connection set for control user i and w_j is peer j's treatment assignment. Correlation between control outcomes and e_i motivates investigation, but graph structure and user differences can also explain it. Comparisons between high- and low-exposure control cohorts therefore require an assignment design and analysis that address those differences. Temporal analysis adds evidence about when effects appear, without by itself distinguishing contagion from a shared external trend. These diagnostics make the interference assumption testable within the observed graph and exposure definition; a null result does not prove that all spillover is absent.

12.5.4.4 Mitigation approaches

When SUTVA violations are detected, four experimental designs modify assignment to address distinct interference paths, each demanding an exposure model tailored to its randomization unit. In graph cluster randomization, treatment is assigned at the community level rather than to individual users. Graph partitioning algorithms like Louvain or spectral clustering seek groups with dense internal connections and sparse cross-cluster edges. Whether that grouping sufficiently limits cross-arm interaction must be measured on the resulting partition. Assignment now operates over

K clusters, so precision depends on the number of independent assignment units, cluster sizes, within-cluster correlation, and remaining spillovers.

Ego-exclusion designs instead filter the analyzed cohort by discarding users whose network exposure crosses a safety threshold. Restricting evaluation to control users with minimal treatment connections (such as $e_i < 0.05$) reduces measured interference, but it cannot guarantee an uncontaminated baseline. Under uniform random assignment, a highly active user with hundreds of connections has a treated-neighbor fraction tightly concentrated around the global treatment share and almost never falls below 0.05, whereas an isolated user with only one or two connections frequently has zero treated peers. Conditioning on low network exposure therefore strips high-degree, high-activity users out of the control cohort and introduces severe selection bias unless the analysis reweights or matches across node degree.

For marketplace dynamics where network connections cannot be partitioned cleanly, switchback experiments alternate treatment and control over time blocks within the selected market or population. This reduces simultaneous cross-arm interaction within that unit, but outcomes can carry over between blocks; the analysis must account for persistence and time trends. Geo-based experiments instead randomize spatial units such as cities, while monitoring cross-boundary travel and shared supply. Choosing the unit trades interference control against the number of independent comparisons and the observation time needed for adequate precision.

The chosen design determines the infrastructure needed to support its claims. Graph-based assignment needs maintained topology and cluster records; a switchback needs an auditable schedule and outcome timing; an exposure-specific estimator needs the peer-assignment information its exposure model uses. Analysis must respect the randomization unit, and monitoring must track the spillovers that could invalidate the comparison. The sensitivity example shows why this work can matter: under its stated assumption, spillover halves the observed lift. The platform must estimate the relevant effect before treating that observed contrast as evidence for or against promotion.

12.5.5 Staged rollout strategies

Cutting 100 percent of live traffic over to a new ranking model all at once exposes every user to a latent bug until operators can roll back, whereas starting at a 5 percent canary cuts that initial blast radius twentyfold. Experiment design fixes how much evidence a release needs; staged rollout fixes how much live traffic may carry the candidate while that evidence accumulates, turning deployment risk into a controllable exposure budget. Traffic increases only after the model earns each larger blast radius through continued acceptable performance.

Napkin Math 12.5: The safety of staged rollouts

Problem: A team is deploying a new ranking model. A blue-green switch that moves all traffic at once (100 percent cutover) exposes all users to any potential bugs. A "Canary" deployment starts at 5 percent traffic. How much does the canary approach reduce the deployment's "Risk Exposure"?

Math: Assume the same request rate, per-request harm, and time to effective rollback in both deployments. Under those assumptions, exposure is proportional to the affected traffic share; comparing shares alone would be insufficient if the smaller canary took longer to reveal the failure.

1. **All-at-Once Blue-Green Exposure**: 100 percent of users affected until rollback.
2. **Canary Exposure**: 5 percent of users affected until rollback.
3. **Risk Mitigation**: $100/5 = 20\times$.

Systems insight: Limiting initial exposure to 5 percent reduces affected traffic by 20× under the shared observation and rollback assumptions. This limits harm while the platform collects live evidence. It does not guarantee detection of rare failures or delayed quality effects, so the release policy must choose its observation window and containment triggers alongside the traffic share.

12.5.5.1 Blue-green deployment

Blue-green deployment keeps the current version (blue) serving while the new version (green) is readied in a parallel environment, then shifts traffic through the routing layer once readiness checks pass. Switching time depends on routing propagation and treatment of in-flight requests. For a large model, keeping both releases ready can reserve substantial accelerator memory. Separate full-capacity fleets duplicate the weights and serving resources during the transition, while a smaller candidate fleet reduces that initial reservation. Fast rollback requires enough ready capacity for the previous version; otherwise model loading and replica startup become part of recovery time. Capacity planning must therefore account for the overlap between releases as well as the steady-state serving allocation.

The environment layout and the traffic policy are separate decisions. A blue-green deployment can switch all traffic at once (the policy notebook 12.5 assumed) or expose the green environment through canary stages. An all-at-once switch forgoes the evidence from limited live exposure, while staged routing retains parallel-environment costs for longer. The release contract determines whether that additional observation and rollback capacity justify their resource cost.

12.5.5.2 Canary deployment

Canary deployment[9] routes a small percentage of traffic to the new version while monitoring for regressions. If metrics remain acceptable, traffic percentage increases until the new version serves all traffic.

A typical progression moves from 1 percent to 5 percent to 25 percent to 50 percent to 100 percent as evidence accumulates. The 5 percent canary in notebook 12.5 is the second of these stages; opening at 1 percent bounds initial exposure further. The duration of each stage matters. Equation 12.7 relates stage duration to sample requirements, request rate, and traffic percentage, giving a lower bound on the time needed to collect the planned sample:

$$T_{\text{stage}} = \frac{n}{\lambda_{\text{arr}} \times p_{\text{stage}}} \tag{12.7}$$

where T_{stage} is the time needed to collect the planned observations, n is the sample requirement from the experiment design, λ_{arr} is the request rate, and p_{stage} is the traffic share expressed as a fraction. This is a collection-time bound, not a complete stopping rule; delayed labels, correlated requests, and the required observation window can extend the stage.

9 **Canary deployment**: Named after the coal-mining practice of using canaries to detect toxic gases. In an ML release, a limited traffic cohort supplies early evidence while bounding initial exposure. Quality regressions can leave health checks green, so the canary needs model and product metrics as well as service-health signals; delayed outcomes can still postpone detection.

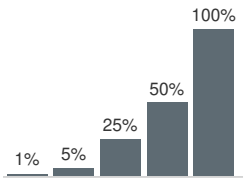

Each canary stage earns a larger share of live traffic.

12.5.5.3 Canary duration calculation

A model serves 1 million requests per hour. Suppose its canary plan specifies approximately 10,000 samples per variant. That is a guardrail sample, sized to catch a large regression quickly before exposure widens, not the sample needed to confirm a small improvement. The sample-size derivation of section 12.5.3.5 takes the baseline rate, effect size, confidence, and power as inputs. Its worked example needs about 30,000 samples per variant to detect a 10 percent relative change at a 5 percent baseline click-through rate. Smaller effects need far more, so promotion evidence comes from the full experiment, while the canary stages below take their guardrail sample as an input.

At 1 percent canary traffic: $T_{\text{stage}} = \frac{10{,}000}{1{,}000{,}000 \times 0.01} = 1$ hour. At 5 percent canary traffic: $T_{\text{stage}} = \frac{10{,}000}{1{,}000{,}000 \times 0.05} = 0.2$ hours $= 12$ minutes.

The organization might configure five rollout stages:

- 1 percent for 2 hours (2× minimum for buffer)
- 5 percent for 30 minutes
- 25 percent for 30 minutes
- 50 percent for 1 hour
- 100 percent deployment

The schedule totals approximately 4 hours. Promotion still requires the planned statistical test and guardrail windows to pass; elapsed time alone does not establish confidence.

Four hours is far shorter than the ensemble soak of table 12.7 or the multi-region waves that section 12.5.7.2 develops later in this chapter, because each schedule answers a different question.

A guardrail sample of fast outcomes, such as errors and clicks, fills in hours. An ensemble change whose effects surface through delayed outcomes needs days of soak. A multi-region rollout repeats evidence collection region by region. The slowest outcome the release must observe, not the traffic share, sets the length of the schedule.

12.5.6 Rollout risk management

Observation windows set how long each rollout stage must run, but a minor bug fix and a core architecture replacement should not share the same exposure schedule or rely on manual intervention when metrics degrade: on a high-throughput ranking service, even a fraction-of-a-point click-through regression compounds into substantial revenue loss during the hours a team spends debating whether an alert is real. Because deployments differ in risk, effective CI/CD systems estimate each change's risk before deciding how quickly to expose it. Equation 12.8 formalizes deployment risk as the product of regression probability, impact severity, and exposure level, providing a quantitative foundation for risk-based rollout decisions:

$$\text{Risk}_{\text{rollout}} = \Pr(\text{regression}) \times \text{Impact}_{\text{regression}} \times E_{\text{exposure}} \tag{12.8}$$

where $\Pr(\text{regression})$ is the probability that the change causes a regression, $\text{Impact}_{\text{regression}}$ is the conditional harm per exposure unit, and E_{exposure} is the number of those units exposed during the rollout period. For example, harm per affected request multiplied by exposed requests gives a total harm estimate. If impact already denotes total loss over the rollout, multiplying by exposure again would double-count it.

The rollout risk framework distinguishes three operational levers. Pre-release evaluation screens out known defects and provides evidence about regression probability ($\Pr(\text{regression})$). Validated fallbacks can reduce conditional harm ($\text{Impact}_{\text{regression}}$) when a defect occurs. Traffic limits, regional isolation, and prompt containment bound exposure (E_{exposure}). Longer observation at a small traffic share can reveal delayed failures, but also exposes that cohort for longer. The release policy must balance evidence collection against accumulated harm rather than assuming that every longer hold reduces risk. Exposure accumulates in time as well as in traffic share. A regression that no alarm catches keeps charging its harm for every hour it stays live, and notebook 12.6 prices one such day.

> Napkin Math 12.6: The cost of a delayed alert
>
> **Problem**: A team deploys a recommendation model that has a silent bug: it reduces click-through rate (CTR) by 0.5 percentage points (10 percent relative), from 5 percent to 4.5 percent. If the service handles 5,000 requests/s and each click is worth \$0.50/click, how much revenue is lost if detection and remediation take 24 hours?
>
> **Math**:
>
> 1. **Total Requests**: 5,000 requests/s × 86,400 s/day = 432,000,000 requests.
> 2. **Lost Clicks**: 432,000,000 requests × 0.005 = 2,160,000 lost clicks.
> 3. **Total Financial Loss**: 2,160,000 lost clicks × \$0.50/click = \$1,080,000.
>
> **Systems insight**: At the assumed traffic and value per click, a 0.5 percentage-point regression costs \$1,080,000 per day. Earlier detection creates value when it enables effective containment sooner. Monitoring and rollback therefore share one operating objective: reducing the time users remain exposed to a harmful release. A drift alarm alone neither establishes the revenue loss nor stops it.

12.5.6.1 Risk categories

The risk equation becomes actionable when its estimates inform a rollout policy. Qualitative examples of that mapping range from a fast canary for a bounded change to additional validation and review for a consequential one (table 12.13). A change label alone does not supply a calibrated regression probability; the decision also needs dependency scope, available evidence, and the consequences of failed containment.

Table 12.13: Risk-Based Rollout Strategy Selection: Four qualitative risk categories mapped to deployment strategies. Teams compute numeric rollout risk as the product of regression probability, impact, and exposure in equation 12.8, then choose the rollout pattern that matches the resulting risk profile.

Category	Typical change (drives Pr(regression))	Typical impact ($\text{Impact}_{\text{regression}}$)	Rollout strategy
Low	Minor code fix	Limited user impact	Fast canary
Medium	Retrained model	Engagement effects	Standard canary
High	New architecture	Revenue impact	Extended shadow + slow canary
Critical	Core model change	Safety implications	Shadow + human review + staged

12.5.6.2 Automated rollback triggers

Risk categories decide how cautiously traffic should expand, but rollback triggers decide when expansion must stop. The trigger configuration in listing 12.5 makes that stop condition executable by binding each monitored metric to a threshold, observation window, and minimum sample count.

Listing 12.5: Automated Rollback Configuration: Metric-specific thresholds, observation windows, and minimum sample sizes that balance sensitivity against false triggers.

```
rollback_config = {
    "metrics": {
        "engagement_rate": {
            "threshold": -0.02,  # 2% relative decline triggers rollback
            "window_minutes": 15,
            "min_samples": 1000,
        },
        "error_rate": {
            "threshold": 0.01,  # 1% absolute increase triggers rollback
            "window_minutes": 5,
            "min_samples": 500,
        },
        "latency_p99": {
            "threshold": 1.5,  # 50% relative increase triggers rollback
            "window_minutes": 5,
            "min_samples": 100,
        },
    },
    "rollback_action": "immediate",  # or 'gradual' for less severe issues
    "notification": ["oncall", "model-owner"],
}
```

Automated rollback must balance sensitivity against false triggers. Minimum sample counts and window durations limit unstable decisions, but they do not define a statistical test. The release contract must also specify the comparison, uncertainty rule, and repeated-monitoring policy. Immediate safety guardrails may require containment before statistical significance is available. These triggers assume a single serving domain. A fleet that spans regions must also keep its version boundary coherent while traffic expands or retracts, which turns rollout and rollback into a coordination protocol.

12.5.7 Multi-region deployment coordination

When Region B switches traffic to a new ranking model before its updated embedding tables finish replicating, or when a regional rollback leaves mobile users bouncing between incompatible model versions across regions, a rollout that passed every single-region check fails globally. A regional rollout must move each model version together with the embedding tables and feature versions it depends on. Where replicas live and how global traffic reaches them is the serving design of section 10.14; operations inherits that placement and must change versions across it safely.

Geographic distribution therefore adds artifact propagation, routing, and regional recovery decisions to the release protocol. Multiple versions may coexist during a staged rollout, but their compatibility, traffic assignment, and rollback behavior must remain within the release contract.

↩ Geographic replication of serving replicas, with the spare regional capacity and gradual traffic return it requires, is established in section 7.11.2.2.

12.5.7.1 Coordination challenges

Multi-region ML deployment couples artifact readiness with input readiness. Large model artifacts consume transfer and loading capacity before a region can serve the new version. For models that retrieve features, receiving the binary is not enough: the region must also provide the feature versions and freshness that the candidate validated. If Region B loads a new model before its required embeddings arrive, requests may fail or use incompatible inputs. The readiness check must therefore cover the model, its feature dependencies, and serving health as one versioned state.

Uncontrolled version mixing is the failure mode that binds multi-region deployment. Staged releases intentionally serve multiple versions, but the protocol must preserve compatibility and experiment assignment even when clocks, traffic volume, routing, and rollback behavior differ by region.

Clock skew and asynchronous progress make timestamps insufficient to establish rollout state. At a planned 2:00 PM UTC start, Region A may begin its 1 percent canary while Region B is still loading the candidate. The coordinator needs explicit readiness and phase acknowledgments to distinguish a permitted staggered release from a failed transition. Regional version differences are acceptable only when routing preserves the declared compatibility and assignment rules.

Regional traffic variation turns the same percentage into different amounts of evidence. A larger canary share in a low-volume region can still collect fewer eligible observations than a smaller share in a high-volume region. The platform must evaluate per-region sample counts, assignment units, and outcome delays against the experiment plan; traffic percentages alone do not establish statistical power.

Cross-region routing can expose version skew when requests move for latency, load balancing, or failover. If an experiment assigns treatment by user, routing must preserve that assignment across regions or record and handle any deviation according to the experiment plan. Mixed versions alone do not invalidate an experiment; unplanned changes in exposure violate the assumptions on which its interpretation depends.

Rollback is the final consistency test. Reverting one region changes the available version set and can violate assignment or dependency assumptions elsewhere. A regional rollback is valid when the release contract permits that state; otherwise the coordinator must contain exposure across the affected regions. Rollback therefore belongs in the rollout protocol rather than an ad hoc emergency procedure.

12.5.7.2 Deployment strategies

Multi-region coordination trades deployment duration against permitted regional skew. In a sequential regional rollout, each region completes its canary progression before the next begins, limiting early exposure while extending the schedule. A globally coordinated rollout instead requires regions to satisfy shared readiness and activation conditions. Logical ordering does not make network delivery simultaneous; routing must enforce the version boundary while activation propagates. A hybrid rollout permits regional progress within declared phase or version bounds, pausing regions when local evidence is insufficient or further progress would exceed the allowed skew. The appropriate policy follows the service's compatibility, assignment, and availability requirements.

An illustrative rollout across 5 regions uses three waves followed by final validation:

1. Canary region (lowest traffic): Full canary cycle, 24 to 48 hours
2. Early adopter regions (2 regions, 20 percent global traffic): Parallel deployment, 24 to 48 hours
3. Majority regions (2 regions, 70 percent global traffic): Parallel deployment, 24 to 48 hours
4. Final validation: Cross-region consistency check, 12 to 24 hours

The assumed windows total 3.5 to 7 days. Actual promotion still depends on regional evidence and guardrails, so low traffic, delayed outcomes, or a failed check can extend the schedule.

12.5.7.3 Traffic management during transitions

Routing controls enforce different parts of the release contract. Region affinity keeps a user's requests in one region, but does not by itself prevent version changes within that region. Version pinning directs requests to replicas serving the assigned release, including during permitted failover. Restricting cross-region fallback can simplify an experiment's exposure boundary, but sacrifices availability when the selected region fails. The platform chooses and tests these controls together so assignment remains interpretable under both normal routing and recovery.

12.5.7.4 Consistency models for deployment

The choice of activation contract affects both coordination cost and the interpretation of deployment metrics. Multi-region systems choose among three consistency models, each trading coordination cost against permitted divergence (table 12.14). Globally ordered activation establishes a shared transition order, eventual convergence permits temporary regional divergence, and bounded staleness limits that divergence. The application must specify which guarantee its decisions require; a financial use case alone does not prohibit controlled experiments or staged versions.

Table 12.14: Consistency Models for Multi-Region Deployment: Version activation contracts trade coordination cost against permitted divergence. Financial or safety-related use does not by itself require identical versions everywhere: assignment consistency, compatible schemas, auditability, and explicitly authorized experiments may permit staged versions. The decision being served sets the required guarantee.

Model	Guarantee	Use case	Coordination overhead
Strong	Globally ordered version activation	One shared activation contract	High (global synchronization)
Eventual	Regions converge to same version	Content recommendations	Low (independent progression)
Bounded staleness	Regions within k_{skew} versions of each other	Real-time ranking	Medium (version monitoring)

For A/B testing, the serving system must implement the declared assignment and record actual exposure. A user assigned to the candidate should receive that candidate across permitted regional routes unless the experiment explicitly models another exposure policy. Unrecorded fallback to the old version changes the intervention being measured and can bias the comparison. Different versions across control and treatment groups are intentional; they are not a relaxation of the assignment contract.

12.5.7.5 Rollback coordination

Rolling back a multi-region deployment requires careful coordination to prevent oscillation and mixed-version serving. The rollback protocol executes across two ordered phases separated by an acknowledgment barrier (figure 12.10). First, the coordinator halts traffic to the new version globally by broadcasting rollback intent, waiting for every region to acknowledge that it has stopped routing new requests to the rejected version, and allowing in-flight requests to drain. Any region failing to acknowledge before the timeout expires is marked unhealthy. Next, the coordinator restores the old version globally once it confirms that all healthy regions serve only the previous release, re-enabling normal traffic routing, clearing transient deployment state, and resetting the system for a future attempt.

The acknowledgment barrier prevents restoration from completing while reachable regions still route new requests to the rejected version. A timeout alone cannot stop an isolated region from serving; the serving path must enforce a lease, routing exclusion, or equivalent fencing rule. The release contract must specify what happens to in-flight requests and to availability during that interval.

Partial rollback allows rolling back individual regions while others continue. This is appropriate when issues are region-specific (infrastructure problems, regional traffic patterns) rather than model-inherent. The deployment coordinator tracks per-region state and prevents inconsistent global decisions based on partial information.

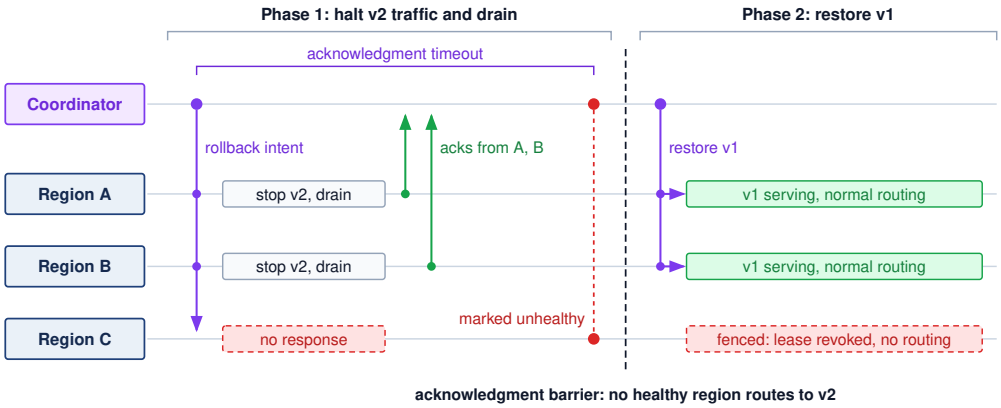

Figure 12.10: Two-Phase Multi-Region Rollback: Time runs left to right on one lane per participant. In phase 1 the coordinator broadcasts rollback intent; regions A and B stop routing new requests to the rejected version (v2), drain in-flight work, and acknowledge. Region C never answers, so when the acknowledgment timeout expires the coordinator marks it unhealthy and fences it. Only after the barrier, when no healthy region still routes to v2, does phase 2 restore the previous version (v1) on the healthy regions; the fenced region stays out of routing until it rejoins under the restored version.

12.5.7.6 Multi-region coordination overhead

In an illustrative five-region deployment with an assumed 80 ms one-way delay per coordination leg, each phase transition needs four synchronous legs (announce intent, collect acknowledgments, confirm completion, collect confirmations), about 320 ms. A canary progression through five traffic stages (1, 5, 25, 50, and 100 percent) adds under two seconds to a rollout measured in hours. Latency is therefore not where coordination costs the release.

The coordination service nevertheless becomes a critical dependency. Its failure policy must specify whether regions may continue serving their last authorized state and whether any further promotion is permitted. A shared activation barrier can halt progress while existing traffic continues; a policy allowing regional independence can permit local progress within its compatibility limits. A bounded-skew contract must stop progress before the bound is exceeded, rather than merely alert afterward. Replicating coordinator state through a consensus protocol, which commits a decision only after a majority of replicas (a quorum) has recorded it, can preserve decisions across replica failures while a quorum remains available. Routing enforcement and regional readiness still determine whether those decisions hold in the serving path.

12.5.8 The end-to-end evaluation contract

Shadow runs, experiments, canary stages, and regional waves each produce release evidence, yet a candidate that passes offline accuracy on one dataset, passes latency benchmarks on a lighter synthetic load profile, and enters an online A/B test without a tested rollback target remains unverified as a whole. Those checks never evaluated the same release state. Evidence from different stages is useful only when it refers to the same candidate and decision. An **evaluation contract** is the versioned release record that binds a model and its dependencies to the workload, metrics, statistical design, acceptance thresholds, rollback conditions, and accountable owners. This record is the subject of the evaluation contract principle (21), one of the fleet-in-production principles listed in the opener to this part of the book. It holds that a fleet release decision is valid only when all of its evidence refers to one versioned contract. The record therefore becomes the unit of validity for everything this chapter has built, from registry versions and feature contracts to rollout stages and experiment statistics. Table 12.15 assembles that record across five evidence planes.

Table 12.15: End-to-End Evaluation Contract: A model release is accepted against one versioned record spanning offline behavior, production-equivalent performance, specialist assurance, online experimentation, and recovery. Each plane has both an evidence owner and a condition that expires its claim.

Evidence plane	Contract must pin	Acceptance or invalidation rule	Accountable owner
Offline behavior	Evaluation-set version, primary metric, consequential slices, baseline, uncertainty evidence, and robustness tests (developed later, in section 14.3.3)	Promote only if prespecified aggregate and slice bounds hold; invalidate when the population, label definition, or intended use leaves the validated envelope	Model owner and domain reviewer
System performance	Hardware and software envelope, open- or closed-loop workload model, concurrency, request mix, warm-up, p50 and tail latency, throughput, errors, and resource cost	Promote only under the production-equivalent load contract; invalidate when capacity headroom or an SLO budget is exhausted	Serving and site-reliability owners
Responsibility and safety	Fairness slices, privacy and safety evidence, affected-party requirements, and assurance claims (developed later, in section 16.2)	Promote only when the named claims have sufficient evidence and review; invalidate on a material change, control failure, or unbounded harm signal	Risk owner and governance authority
Online experiment	Assignment unit, holdout, minimum detectable effect, power, sample size, confidence interval, multiple-testing correction, stopping rule, and guardrail metrics	Promote only after the planned decision rule passes; stop and roll back when any guardrail crosses its prespecified threshold	Experiment owner and product decision owner
Recovery and audit	Rollback artifact and dependencies, trigger windows, evidence-retention period, data and code versions, environment, seeds, raw aggregate counts, and approvals	Reject exposure without a tested rollback; reopen the decision when evidence cannot be reproduced or retained for the required review period	Release owner, platform owner, and audit custodian

Acceptance is deliberately distributed, but promotion authority is not ambiguous. The model and domain owners accept predictive utility; the serving owner accepts capacity and reliability; specialist risk owners accept bounded robustness, fairness, privacy, and safety claims. A named release authority promotes only when all required sign-offs refer to the same artifact. The platform automates gates and preserves evidence, but it cannot choose the product objective or waive a failed assurance claim.

Invalidation is the return path through the contract. Drift beyond the evaluated population, a feature or label-definition change, exhausted latency headroom, a statistically supported guardrail regression, or a material governance change sends the candidate back to the owning gate. Thresholds and rollback targets are fixed before exposure so an apparent primary-metric gain cannot redefine success after the result is known. Reproducibility requires retaining the inputs to that decision, not merely its verdict: the dataset and traffic snapshot, model and code versions, environment and hardware envelope, random seeds, assignment and sampling rules, sufficient statistics or permitted logs, approvals, and the exact rollback artifact.

The fraud shadow workflow of section 12.5.3.2 shows how one stage's evidence maps onto these planes. Its labeled precision and recall, with the paired outcomes, label maturity, intervals, and recall tolerance its gate required, are paired live-traffic quality evidence that the contract records. They do not discharge the online-experiment plane, because a shadow run has no randomized assignment and exposes no user to the candidate's outputs. The workflow's p99 latency is system-performance evidence, but the sampled phase saw only 10 percent of traffic; the full-shadow phase would supply the production-load measurement if its latency were recorded. The clipping change then triggered an invalidation rule. Because it changed the model's inputs, it voided that evidence and sent the revised candidate back through its gates. The workflow as described names no tested rollback artifact, raw aggregate counts, evidence-retention period, or approvals, so its recovery-and-audit plane would still be empty before canary exposure.

12.5.9 CI/CD patterns by model type

Every workload family shares the five evidence planes of table 12.15, yet the intended use determines which plane gates promotion. Releasing a general-purpose LLM with a two-hour automated check

misses subtle safety regressions, whereas holding an emergency fraud-model patch in a two-week human review window lets an active attack drain funds. Table 12.16 returns to the workload families of table 12.2, whose cadence and rollback targets still apply, and adds the column that table lacked, what each family must validate before promotion. The listed durations are operating assumptions, not consequences of architecture; label maturity, traffic, and potential harm determine the actual validation and rollback policy.

Table 12.16: CI/CD Patterns by Model Type: The rows illustrate validation and response policies for different workloads. Their durations are planning assumptions; promotion depends on the required evidence, including outcome-label delay, and rollback depends on a ready compatible target.

Pattern	Model type	Validation focus	Rollout speed	Rollback speed
Quality-gated	LLM	Human eval, safety	Days to weeks	Minutes with warm standby
Metric-driven	Recommendation	Engagement metrics	Hours to days	Minutes
Threshold-gated	Fraud	Precision/recall	Hours	Seconds
Accuracy-focused	Vision	Classification metrics	Days	Minutes

For a broadly capable LLM, evaluation must cover consequential uses that aggregate benchmarks can miss. Automated tasks such as MMLU and HumanEval screen selected capabilities; human review and adversarial probes examine additional failure modes; shadow traffic tests operational behavior on live inputs. A planning window of 2–4 weeks may accommodate these activities and staged observation, but elapsed time is not the acceptance criterion. The required evidence depends on the changed capabilities and the service's exposure policy.

Recommendation pipelines balance freshness with statistically defensible comparison. Offline NDCG and recall can screen candidates, interleaving can compare ranker preference, and A/B tests measure the selected product outcomes. A 24–48-hour rollout is feasible only when traffic and outcome timing supply the required evidence within that interval. Promotion and rollback follow the prespecified rules rather than a fixed refresh deadline.

Fraud models add adversarial urgency, but confirmed outcomes can arrive after the decision window. Historical labeled cases and shadow scoring provide evidence before exposure; neither makes fresh precision and recall immediately observable. The release policy must state what evidence supports an urgent change, what uncertainty remains, and who accepts its bounded exposure. Rapid routing to a validated fallback can contain harm while additional labels accumulate, provided that fallback remains suitable for the current threat.

Edge fleets, absent from the table, apply the same contract to weeks-long rollouts that leave many model versions active at once, to devices whose memory and thermal limits require hardware-in-the-loop validation, and to telemetry that privacy rules limit to federated aggregates. Section 11.10.1.5 develops these release controls, including why the devices that report health do not represent the deployed fleet.

Whether a model runs in a data center or on a phone, a CI/CD pipeline promotes candidates supported by evidence within a specified operating envelope. Production can move outside that envelope as data, traffic, and dependencies change. Monitoring must therefore preserve the release contract after promotion, detect when its evidence no longer applies, and return the affected model to evaluation or recovery.

12.6 Monitoring at Scale

Preserving a release contract after promotion requires knowing whether a production regression comes from shifting data, an upstream dependency, or the hardware underneath the model. Because application-layer drift and latency checks depend on a healthy physical substrate, where a degraded NVLink connection or a thermally throttling GPU can masquerade as a software timeout, fleet monitoring starts bottom-up with hardware telemetry and climbs through alert aggregation, model-quality signals, and dashboards so operators can route each symptom to the layer that can repair it.

12.6.1 Fleet telemetry and hardware observability

Under an illustrative independent-failure model, a 50,000-hour per-GPU planning interval gives a 10,000-GPU cluster a hardware failure roughly every five hours, and the full-system interval of section 7.2.2 is shorter still. Because failure is routine at pod scale, fleet monitoring and hardware observability are prerequisites for productive training. Recovering in 10 minutes rather than hours depends on detection, usable recovery state, replacement capacity, and restart procedures, so the monitoring system is itself a distributed service collecting telemetry across devices, switches, and cooling components.

The telemetry spans three physical levels. At the GPU level, junction temperature tracks cooling degradation, and error-correcting code (ECC) counts in high-bandwidth memory (HBM) and SRAM track memory health. A rising rate of correctable errors can precede an uncorrectable, job-ending error. Streaming multiprocessor (SM) utilization shows whether the device is doing work. None alone identifies the cause: low utilization can reflect input starvation, synchronization, or a device problem. At the network level, retry counters on individual ports can point to a failing cable or switch; at the facility level, coolant temperature deviations can trigger load shedding before thermal limits force throttling. Correlating these signals with workload and topology lets responders narrow the failure domain before deciding whether to reduce load, quarantine hardware, or investigate software.

Correlation analysis helps locate the scope of a shared disturbance. Elevated temperature on one GPU motivates local inspection; simultaneous changes across 8 GPUs in a node suggest checking shared cooling and workload changes. A rack-wide pattern adds the coolant distribution unit (CDU) and rack-level load to the investigation, while a multi-rack pattern raises facility-level hypotheses. These patterns guide diagnosis rather than prove a cause. Grouping related alerts preserves that context and avoids presenting one shared event as hundreds of unrelated failures.

Telemetry capacity follows the series count, sample rate, bytes per observation, retention, and indexing policy. Sampling one metric every second across 10,000 GPUs produces 864 million observations per day before other metrics or labels are added. A daily byte total or monitoring-budget share cannot be inferred without the encoding and retention assumptions. The platform therefore bounds label cardinality and provisions ingestion, storage, and queries as an explicit platform workload.

Before allocating a long training job, a scheduler can use automated preflight checks to screen its nodes. General matrix multiply (GEMM) benchmarks exercise arithmetic throughput, NCCL AllReduce tests exercise collective paths, and memory stress tests look for errors under load. Results must be compared with baselines for the same hardware and configuration before a node is quarantined. Preflight consumes startup time and cannot reveal every fault; its value depends on the failed work it avoids and the cost of unnecessary exclusions.

A gray failure leaves a component responsive while degrading its useful performance. In synchronous data-parallel training, a persistently slow participant can delay collective completion even when every device passes an availability check. Per-rank timing and hardware telemetry help localize the slowdown; controlled microbenchmarks can then compare suspected nodes with healthy peers. An illustrative policy flags GEMM throughput below 90 percent of a matched fleet median or NVLink bandwidth below 85 percent. Those thresholds require calibration for hardware generation, software, load, and measurement variability. A low result warrants investigation rather than proving a hardware defect, because configuration and contention can produce similar symptoms.

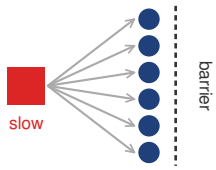

One slow accelerator stalls every peer at the barrier.

In multi-tenant clusters, shared links can also couple otherwise separate workloads, so fabric telemetry must attribute congestion to the tenant that caused it. The isolation mechanisms that bound such interference, along with their limits, are established in section 3.8.2.

Checkpoint cadence is the point where hardware telemetry becomes visible to ML operators rather than only to data-center technicians. A degraded node that forces frequent checkpoint recovery, or a storage path that lengthens checkpoint writes, turns directly into lost training throughput. Section 7.8.1 derives the checkpoint-compute trade-off and the Young-Daly interval; the operations responsibility here is to expose the telemetry those formulas depend on: write time, failure rate, straggler behavior, and recovery duration. Without that visibility, a platform may report that a job is "running" while useful learning has quietly collapsed behind retries and slow checkpoints.

↩ Gray failures and the straggler detection that exposes them are established in section 7.9.5.

With the physical substrate accounted for, monitoring returns to the deployed model fleet. Validation and canary evidence cover a candidate under observed conditions; production can later change abruptly or drift over weeks and months. Monitoring must therefore retain acute-failure controls while detecting gradual loss of quality. Independently paging on every model metric can multiply notifications from a shared event and create alert fatigue. Fleet-level aggregation connects those symptoms while preserving urgent local signals that require immediate action.

12.6.2 The alert fatigue problem

Joining hardware counters with per-model quality, drift, and latency metrics across a portfolio multiplies the number of time series evaluated on every check interval. Even when each metric uses a well-calibrated statistical threshold with a tiny per-check false-positive rate, evaluating thousands of checks every few minutes produces a continuous stream of random threshold crossings whose operational cost depends on the alerting policy.

> Napkin Math 12.7: False alarms across many checks
>
> **Problem**: Consider a system monitoring 100 models. Each model has 10 metrics (latency, accuracy, drift, and others). Alert thresholds are set at 3-sigma (99.7 percent specificity), and a control script re-evaluates every metric on a fixed interval. Assume calibrated stable Gaussian null behavior and count raw threshold crossings before grouping or suppression; those crossings are not necessarily distinct incidents.
>
> **Math**:
>
> 1. **Total Monitors**: 100 models $\times$ 10 metrics = 1,000 monitors.
> 2. **False Positive Rate**: $1 - 0.997 = 0.003$ (0.3 percent).
> 3. **Checks per day**: Assume checks every 5 minutes (288 checks/day).
> 4. **Daily false threshold crossings**: $1{,}000 \times 288 \times 0.003 \approx 864$ alerts/day.
>
> **Systems insight**: Even with high-specificity (3-sigma) alerts, scale becomes problematic. Aggregation, deduplication, persistence windows, and calibrated thresholds reduce repeated notifications. Specific safety or hardware-failure alerts must still trigger independently when their contracts require immediate action.

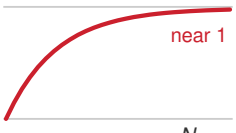

False alerts become near-certain as checks multiply.

The multiple-testing problem from concurrent experiments (section 12.5.3.7) also appears in monitoring. For independent checks under the null, each with false positive rate α_{fp}, equation 12.9 shows that the probability of at least one false alert approaches certainty as the number of tests N_{tests} grows, because the probability of no false alert decays exponentially:

$$\Pr(\text{at least one false alert}) = 1 - (1 - \alpha_{\text{fp}})^{N_{\text{tests}}} \tag{12.9}$$

The equation holds at any per-check rate, and teams under pressure to catch regressions sooner often loosen thresholds. Loosening the notebook's 3-sigma thresholds to 2 sigma raises the per-check rate from 0.3 percent to roughly $\alpha_{\text{fp}} = 0.05$. With that rate and $N_{\text{tests}} = 1000$ (100 models $\times$ 10 metrics):

$$\Pr(\text{false alert}) = 1 - (1 - 0.05)^{1000} = 1 - 0.95^{1000} \approx 1.0$$

The probability is essentially 100 percent for this collection of independent checks. This result establishes the likelihood of at least one false crossing, not the fraction of alerts that are false. Repeated checks compound the operational burden when every crossing generates a page, and frequent nonactionable pages can train operators to discount the alert stream, making genuine incidents harder to recognize. For the notebook's fleet of 100 models with 10 metrics checked every 5 minutes, the 5 percent rate raises the daily raw crossings far above the 3-sigma count.

Expected daily false alerts: Daily false alerts $= 100 \times 10 \times 0.05 \times \frac{24 \times 60}{5} = 14{,}400$.

Even if 99 percent of these crossings are deduplicated or auto-resolved, the remaining 144 alerts daily would impose a substantial interruption burden if each became a page. Comprehensive

telemetry therefore needs a policy that separates diagnostic records from notifications requiring immediate action.

12.6.3 Hierarchical monitoring architecture

When a shared feature store stalls, paging twenty downstream model teams simultaneously floods on-call rotations without telling anyone that the root cause sits in shared infrastructure. To prevent uncoordinated alarm floods, hierarchical monitoring arranges telemetry as a four-tier triage pyramid that maps operational scope from component health at the base to fleet-wide business impact at the apex (figure 12.11). Broad signals reveal fleet impairment, portfolio signals route incidents to the owning model family, and granular model or infrastructure telemetry supports local diagnosis. Routine fluctuations remain diagnostic records, while urgent local failures still page their owners without waiting for a fleet-wide signal. The hierarchy provides context for routing and grouping alerts while preserving those immediate response paths. It routes evidence to the owner of the affected level and does not decide the remedy; whether a model must change remains its owner's decision.

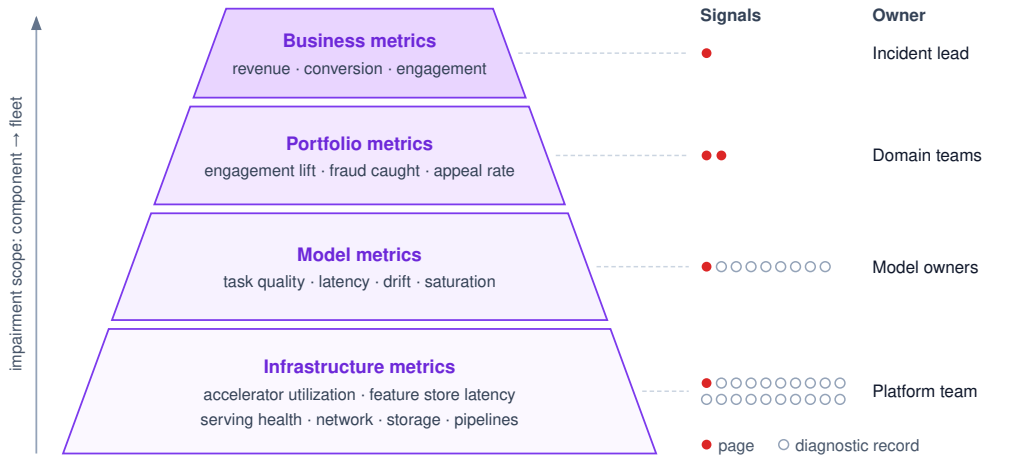

Figure 12.11: Hierarchical Monitoring Architecture: To prevent alert fatigue, monitoring operates at four levels, each with its own owner. Business and portfolio metrics raise the few alarms that indicate broad impairment and route each incident to the owning model family, while model and infrastructure metrics serve primarily for diagnosis and root-cause analysis.

The hierarchy works because each level answers a different routing question and has a different owner, from the incident lead who watches business impairment to the platform team that repairs capacity, networking, and storage failures (table 12.17).

Table 12.17: Hierarchical Monitoring Levels: Monitoring levels organize incident scope and ownership. Broad signals route shared incidents, while urgent local failures can page their owners directly; evidence at every level supports diagnosis of model, data, serving, or infrastructure causes.

Level	Representative signals	Primary owner	Operational role
Business	Revenue or conversion attributed to recommendations, engagement indicators, automation rate, and human-review volume.	Incident lead	Few high-confidence alerts that indicate product or business impairment.
Portfolio	Engagement lift and diversity for recommendation, fraud caught and false-positive rate, or violation detection and appeal rate.	Domain team	Routes investigation to the model family or product area that owns the shared objective.
Model	Task quality, latency distributions, throughput, error rates, resource utilization, serving saturation, and recent deployment state.	Model owner	Supports local diagnosis after business or portfolio signals move.

Level	Representative signals	Primary owner	Operational role
Infrastructure	GPU cluster utilization and availability, feature store latency, training pipeline time, serving health, networking, and storage.	Platform team	Routes failures that require capacity, placement, networking, storage, or control-plane repair.

12.6.3.1 Multidimensional reliability budgets

A fleet cannot summarize reliability with availability alone. An endpoint may return HTTP 200 while serving stale features, missing its latency target, or producing predictions below the quality floor. Each service therefore needs an SLO vector

$$\mathbf{S} = (S_{\text{availability}}, S_{\text{latency}}, S_{\text{freshness}}, S_{\text{quality}}),$$

with one indicator, target, measurement window, and owner for every dimension. For dimension i, the normalized budget consumption over a window is

$$B_i = \frac{V_i}{V_i^{\text{allow}}},$$

where V_i is the fraction of eligible observations that violate target i and V_i^{allow} is the allowed violating fraction for that same target and window. The numerator and denominator have identical units within each dimension: failed-request fraction for availability, late-request fraction for latency, stale-prediction fraction for freshness, and below-floor evaluation fraction for model quality. The quality contract must define its evaluation units, label delay, and coverage; unlabeled predictions cannot silently count as passing observations. A value $B_i = 1$ exhausts that dimension's budget; $B_i > 1$ exceeds it. Ratios may be compared after this normalization, but raw latency, freshness, and quality magnitudes are not commensurable. The fleet remains inside its measured reliability envelope only when $\max_i B_i \leq 1$, the reliability condition of the evaluation contract (principle 21). Averaging would hide a severe regression behind healthy dimensions. A service whose freshness budget is overspent can still report a comfortable mean consumption, yet the maximum stops its promotion (figure 12.12). Missing or stale evidence requires its own release policy rather than an assumed pass.

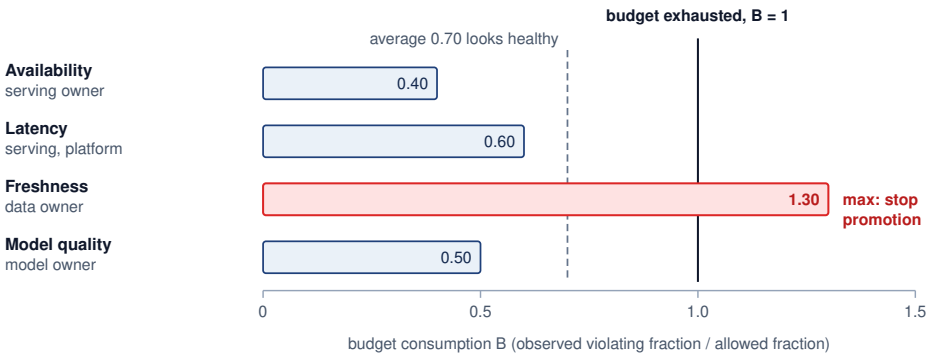

Figure 12.12: Reliability as a Vector of Budgets: Each bar is one dimension's budget consumption, its observed violating fraction divided by the allowed fraction, with the owner of that dimension beside it. The values are illustrative. Availability, latency, and quality sit inside their budgets and the four-dimension average is 0.70, but freshness has spent 1.30 of its budget; because the maximum, not the average, governs, the service stops promotion until the data owner restores freshness.

Burn rate converts this vector into an operating control. For dimension i, the **burn rate** is the observed violating fraction over a short recent interval divided by the allowed fraction V_i^{allow}, so it measures how fast the budget is being spent relative to the pace that would exhaust it exactly at the end of the window. A burn rate of 1 spends the budget on schedule; a burn rate of 10 would exhaust the whole window's budget in a tenth of the window. A fast burn in any dimension pages the

owning team and stops promotion. Sustained moderate burn reduces rollout velocity and requires remediation. An exhausted budget freezes nonessential releases until the responsible indicator returns inside its recovery policy.

The platform owns indicator computation, common windows, alert routing, and enforcement. Model teams own task-quality thresholds and remediation for model behavior; data owners own freshness and validity; serving owners own latency and availability. This division turns monitoring into an explicit contract between evidence and authority rather than a dashboard that everyone can see but no one must act on. Each exhausted dimension triggers a release consequence its own owner carries out (table 12.18).

Table 12.18: ML Reliability Budget Vector: Fleet reliability joins service health, data freshness, and model behavior. Any exhausted dimension can stop a release because success in one dimension cannot compensate for failure in another.

Budget dimension	Representative indicator	Primary owner	Release consequence
Availability	Successful eligible requests/total requests	Serving	Halt rollout or fail over when fast burn exceeds policy
Latency	Requests meeting the tail-latency objective	Serving/platform	Reduce traffic, capacity pressure, or candidate complexity
Freshness	Predictions using features inside freshness SLO	Data	Block promotion and fall back to a safe feature path
Model quality	Slice-aware task metric above its quality floor	Model/domain	Roll back, abstain, or route to a validated baseline

12.6.4 Anomaly detection across the fleet

Protecting those multidimensional reliability budgets requires detecting correlated failures and input shifts before a single upstream fault exhausts every downstream model's error allowance. When 60 of 100 models cross their control limits at the same timestamp because a shared feature pipeline broke, or when user input distributions shift days before ground-truth labels arrive, per-model static thresholds cannot distinguish a local bug from a fleet-wide event. Fleet-wide anomaly detection adds portfolio context to individual metric thresholds. The detection stack proceeds from local statistical process control, to correlation across model families, to drift scores that characterize changes in data distributions. Each stage narrows the investigation, but none establishes a cause by itself. Grouping related signals can reduce duplicate pages while retaining the evidence needed to route an incident to the right owner.

12.6.4.1 Statistical process control

Statistical process control originated in industrial quality control (Shewhart 1931). When adapted for ML monitoring, control charts[10] track whether metric distributions remain stable over time and complement drift-detection methods that watch for changing data distributions (Gama et al. 2014). The core idea is distinguishing common cause variation (normal fluctuation) from special cause variation (genuine anomalies).

For a metric X with established mean μ and standard deviation σ, the upper control limit is $\text{UCL} = \mu + 3\sigma$ and the lower control limit is $\text{LCL} = \mu - 3\sigma$. Points outside control limits or systematic patterns, such as 7 consecutive points above or below the mean, trigger investigation. Adding pattern rules changes the combined false-alarm rate; the single-check 3-sigma rate no longer describes the entire policy. Calibration must also account for seasonality and temporal dependence in the monitored metric.

12.6.4.2 Fleet-wide correlation

When multiple models show similar anomalies simultaneously, shared infrastructure or an upstream data pipeline becomes a candidate cause. Correlation analysis across the fleet supports three operational tasks: identifying shared dependencies to investigate, grouping related alerts into an incident, and prioritizing triage based on the breadth of downstream impact. As shown in listing 12.6, the fraction of anomalous models can guide escalation from model-by-model debugging to a coordinated investigation. Crossing that threshold does not establish a shared failure; dependency records, deployment timing, and component telemetry must support the attribution.

[10] **Statistical process control (SPC)**: Shewhart (1931) established the control-chart framework for manufacturing quality control. Under a normal approximation, a 3-sigma control limit produces a false-alarm probability of only 0.27 percent for a stable process. For ML fleet monitoring, even small per-test false-alarm rates compound across many models and metrics, making alert-fatigue management an inherent constraint of applying SPC at platform scale.

Listing 12.6: Fleet Anomaly Attribution: Detecting correlated anomalies across a model fleet and attributing them to shared infrastructure or data causes.

```
def detect_fleet_anomaly(model_metrics, threshold=0.6):
    """
    Detect correlated anomalies across model fleet.

    Returns list of (timestamp, affected_models, likely_cause) tuples.
    """
    anomalies = []

    for timestamp in model_metrics.timestamps:
        # Identify models with anomalous metrics at this time
        anomalous_models = []
        for model in model_metrics.models:
            if is_anomalous(model_metrics[model][timestamp]):
                anomalous_models.append(model)

        # Check if anomaly fraction exceeds correlation threshold
        if (
            len(anomalous_models) / len(model_metrics.models)
            > threshold
        ):
            # Many models affected -> likely shared cause
            cause = attribute_to_shared_cause(
                timestamp, anomalous_models
            )
            anomalies.append((timestamp, anomalous_models, cause))

    return anomalies
```

The threshold controls when the monitor escalates to a fleet-wide investigation. A low threshold increases sensitivity and investigation load; a high threshold can miss failures affecting a smaller but critical subset of models. Above the threshold, the incident coordinator uses shared dependencies to identify candidate owners, while affected model teams retain responsibility for local containment.

12.6.4.3 Drift detection

Distribution changes can invalidate the population assumptions behind a release even while service health remains normal. Monitoring must distinguish a change in inputs from a change in the input-output relationship; neither an input alarm nor a healthy input histogram alone establishes model quality.

Covariate shift moves the input distribution $p(x)$ while the input-output relationship holds, so label-free input monitoring can detect the move but cannot confirm that the relationship stayed fixed. Concept drift changes the relationship $p(y \mid x)$ itself, so detecting it requires ground-truth labels that often arrive minutes, days, or weeks late.

Because labels are often delayed, most real-time monitoring systems use input drift as a leading indicator of possible performance degradation. A cheap drift score is valuable even when it is not a complete theory of robustness because it tells the operator where to investigate before outcome labels arrive.

For continuous features, the population stability index (PSI)[11] quantifies distribution shift (Yurdakul and Naranjo 2020). Equation 12.10 computes PSI as the sum of log-ratio weighted differences between actual and expected bucket proportions. In platform operations, PSI plays a narrow role as an inexpensive alerting signal.

$$\text{PSI} = \sum_{b=1}^{n_{\text{bins}}} (A_b - E_b) \times \ln\left(\frac{A_b}{E_b}\right) \tag{12.10}$$

where A_b is the proportion in bucket b of the actual (current) distribution, E_b is the proportion in bucket b of the expected (reference) distribution, and n_{bins} is the number of buckets.

An illustrative policy uses these coarse bands to route investigation. PSI below 0.1 leaves the model in the normal monitoring path; values between 0.1 and 0.25 create an investigation ticket;

↪ The full shift taxonomy, including label shift, and the evidence that distinguishes one kind of drift from another are developed in section 14.4.1.

[11] **Population stability index (PSI)**: Bands such as PSI < 0.1 for routine monitoring and PSI ≥ 0.25 for investigation are heuristics, not guarantees of stability or model failure. Once histograms are available, the comparison uses bucket proportions rather than individual records. Histogram construction, sample size, binning, and treatment of empty buckets still affect cost and interpretation, so the platform must calibrate the signal on its own data.

values at or above 0.25 escalate review of the shifted inputs and affected predictions. These actions depend on local calibration. A high PSI does not establish degraded model quality, and a low PSI does not establish that quality remains acceptable.

Napkin Math 12.8: Time-to-detection: Monitoring lag

Problem: A production model has a baseline accuracy of 95 percent. A data drift event causes accuracy to fall to 93 percent. If labeled outcomes arrive at 1000 samples/hour, what collection time supports the stated detection-power calculation, assuming an independent baseline sample of the same size and immediately available labels?

Math: Detection requires enough samples to distinguish the signal (the 2-percentage-point drop) from the noise (random variance). The sizing logic of equation 12.4 applies, in its unpooled form because the baseline and degraded accuracies (a_1, a_2) differ, with the same confidence and power targets ($z_{\alpha/2}$, z_β):

$$n \approx \frac{(z_{\alpha/2} + z_\beta)^2 \left[a_1(1-a_1) + a_2(1-a_2)\right]}{(a_1 - a_2)^2}$$

1. **Samples Required**: At 95 percent confidence and 80 percent power, the standard normal quantiles are $z_{\alpha/2} = 1.96,\ z_\beta = 0.84$, and plugging in gives $n \approx (1.96 + 0.84)^2 \times [0.95(1-0.95) + 0.93(1-0.93)]/(0.95-0.93)^2 \approx 2,207$ samples.
2. **Detection Latency**: 2,207 samples / 1000 samples/hour $\approx$ 2.2 hours.

Systems insight: Monitoring is not instantaneous. The computed duration is a planned collection time at 80 percent power, not a guaranteed detection deadline. It assumes an independent baseline sample, the stated post-change rate, and immediately available labels. Delayed labels, dependence, and repeated testing change the detection contract. Input drift can provide an earlier investigation signal, but it does not establish a quality regression by itself.

Fleet-wide drift monitoring extends that signal across the portfolio. A PSI spike in one noncritical feature may be a local investigation, while simultaneous drift across shared features or dependent models warrants a broader investigation. A shared data pipeline failure is one candidate cause; a genuine change in the population served can produce the same pattern. Pipeline checks and outcome evidence distinguish a data defect from a change to which the models must adapt. Whether a confirmed shift calls for retraining, adaptation, or no change is a model decision developed in section 14.4.2.5; the fleet monitor's responsibility ends when calibrated evidence reaches the owner who makes that decision.

12.6.5 Model-type specific monitoring

A fraud detector watching for active card-testing attacks, a recommendation ranker watching real-time click-through rates, and an LLM assistant watching for toxic or hallucinated generations cannot share one monitoring metric or polling interval. Monitoring cadence follows failure cost, label delay, and traffic volume. The same workload families that set release cadence in table 12.2 and validation focus in table 12.16 now need a post-promotion counterpart: which metric each watches and how often. The illustrative parameters in table 12.19 must be calibrated to the service, because a frequently updated metric is useful only if enough observations support the alert decision.

Table 12.19: Illustrative Monitoring Policies: Example metrics, thresholds, and collection frequencies. None is a universal requirement of a model family. Label availability, sample size, failure cost, and the required detection delay determine the production policy.

Model type	Primary metrics	Alert thresholds	Monitoring frequency
Recommendation	CTR, engagement lift	5% relative drop	Real-time
Fraud Detection	Precision, recall, fraud rate	1% degradation	Real-time

Model type	Primary metrics	Alert thresholds	Monitoring frequency
LLM	Quality scores, safety metrics	Per-model calibration	Hourly
Vision	Accuracy by class	Dataset-specific	Daily
Search Ranking	NDCG, click position	2% degradation	Real-time

Recommendation monitoring must respond quickly to user impact, but each metric becomes informative on its own observation timescale. Click-through rate, dwell time, and conversion rate should be compared against time-matched historical baselines, control traffic when available, and the previous model version after a release. Conversions can arrive later than clicks, and sparse traffic can delay a reliable comparison. Engagement metrics alone also leave gaps; diversity, catalog coverage, and filter-bubble indicators track longer-term concerns, while revenue attribution and promotional-inventory metrics connect recommendation quality to business outcomes.

Fraud monitoring binds on adversarial urgency. Missed fraud creates immediate financial loss, but excessive false positives create customer friction and manual-review load. The monitoring surface therefore has to pair detection rate, prevented dollar amount, and detection latency with false-positive rate, blocked-legitimate-transaction events, manual-review volume, and adversarial indicators such as probing patterns or sudden shifts in fraudulent behavior.

LLM monitoring binds on semantic quality, which is difficult to measure from service metrics alone. Latency, token generation rate, error rate, and safety-classifier scores establish operational health; user satisfaction signals such as thumbs-up/thumbs-down rates, regeneration rate, and task-completion proxies add delayed quality evidence. Safety metrics such as toxicity detections, refusal rate, and hallucination indicators remain partial, so they should trigger review rather than claim complete coverage.

Service-health metrics alone cannot detect semantic safety failures such as persuasive misinformation or skilled manipulation. Here, red teaming supplies an additional monitoring channel: human evaluators or automated probes attempt jailbreaks before deployment, and a smaller continuous probe set checks production behavior against known failure cases. Sampling outputs for delayed human review adds evidence about failures that automated metrics miss. Neither a passing probe set nor a reviewed sample establishes complete safety coverage.

12.6.6 Observability architecture

When a single user request traverses ten inference services and suffers a 50 ms tail-latency spike or a silent accuracy regression, independent service counters cannot tell investigators which stage stalled. Unsampled logs from 10,000 nodes can flood the telemetry backend during the outage. Effective monitoring therefore requires observability infrastructure that preserves enough context to route an alert from symptom to owner. The metrics, logs, and traces that make up that infrastructure were introduced for debugging a single distributed job; at fleet scale the same channels must route evidence across many models and owners. Metrics show that a service property moved outside its envelope, traces show where a request spent time, logs identify which component emitted the abnormal event, and prediction logs connect system health back to model behavior. Distributed tracing,[12] exemplified by Google's Dapper system (Sigelman et al. 2010), preserves a request's path across inference services so investigators can locate latency contributions. Taken together, the channels form an evidence map rather than a set of disconnected telemetry tools, with each channel matched to the operational question it answers (table 12.20). The channels join into one map only when their records share identifiers, such as the request, the model version, and the host, so that a metric, a trace, and a log from the same incident can be matched. Section 12.7.5 specifies that identifier contract for distributed jobs.

[12] **Distributed tracing**: Google's 2010 Dapper paper describes propagating trace context across service boundaries while controlling collection overhead through sampling and instrumentation in common libraries. In a multi-model pipeline, the resulting trace connects stage timings to the same request rather than requiring investigators to infer the path from independent service summaries. Without it, a 50 ms tail-latency spike in a 10-model pipeline must be investigated in all 10 models separately. Sampling and instrumentation coverage determine which requests and stages remain visible.

Table 12.20: Observability Signals by Operational Question: Observability architecture is an evidence-routing system. Metrics trigger attention, traces localize the request path, logs explain component behavior, and prediction logs connect operational symptoms to model quality.

Evidence channel	Primary question	What it preserves	Typical failure revealed
Metrics	Which service property moved out of bounds?	Real-time streams for alerting and aggregated time series for trends	Drift, saturation, error-rate spikes, cost anomalies
Traces	Where did this request spend time?	Request IDs propagated across model invocations, timing, and resources	Cross-model latency, fan-out failures, hot stages
Logs	Which component emitted the abnormal event?	Structured events with consistent schemas and searchable indexes	New error classes, dependency failures, bad releases
Prediction logs	Did system health preserve model behavior?	Sampled inputs, outputs, features, and delayed labels	Accuracy regressions, bad cohorts, data gaps

↩ The observability pillars for a single distributed job are established in section 7.13.2.

The channels also grow at different rates, which decides what the platform can afford to collect continuously. Metrics stay compact only while the number of distinct time series is bounded. Each label combination creates another series, so a request identifier or an unbounded user label can overwhelm monitoring even when each sample is small. Logs grow with recorded events and traces with spans, so both are sampled or aggregated at fleet scale. When 10,000 edge nodes or hundreds of microservices emit at once, unsampled telemetry can become a thundering herd, a simultaneous burst of writes that perturbs the production system it is meant to observe. Effective platforms therefore present high-level metric dashboards by default and open lower levels only when an anomaly warrants investigation.

Prediction logging can be expensive because it records model-facing data rather than compact service counters. Production systems therefore sample it deliberately, for example logging 1 percent of predictions or retaining additional examples from cohorts under investigation. The sample must serve the intended analysis; a debugging sample concentrated on failures does not directly estimate fleet-wide accuracy. Recording the selection policy alongside inputs, outputs, and eventual labels preserves the context needed to interpret the evidence without making every prediction a permanent storage record.

12.6.7 Dashboard design

Sampling metrics, logs, and traces keeps telemetry storage bounded, yet a single all-purpose dashboard still fails both ends of an outage. An executive checking whether the business is impaired gets lost in per-GPU trace spans, while an on-call responder diagnosing a tail-latency spike is stuck staring at a high-level portfolio summary. A dashboard is the human attention surface of the monitoring hierarchy, and it should answer routing questions without forcing every reader into the same debugging console. Executives need to know whether the business is impaired, domain owners need to know which portfolio is failing, model owners need the evidence for their model, and incident responders need the trace from symptom to root cause. Each view therefore answers one question at one level of detail (table 12.21).

Table 12.21: Dashboard Views by Operational Question: Dashboard hierarchy should route attention from business impact to root cause. Each view reveals the amount of detail needed for the decision at that level, preventing executive views from becoming debugging consoles and investigation views from becoming portfolio summaries.

View	Question answered	Signals to surface
Executive	Is the business impaired?	Platform health, business impact, active incidents, key trends
Portfolio	Which domain or model family is driving the impairment?	Model inventory, portfolio metrics, recent deployments, resource utilization and cost
Model	What changed for the responsible model?	Current metrics vs. baselines, deployment history, drift indicators, cost attribution
Investigation	Why did the change occur?	Cross-model correlations, time-series overlays, log search, request-level trace detail

The same views carry cost as one more monitored dimension: the portfolio and model views surface spend beside quality and latency, and section 12.9.7.3 develops the anomaly detection that watches that spend. Even a well-instrumented fleet does not prevent failures. The monitoring hierarchy detects them; what remains is the response, which must classify the symptom, contain its impact, and route it to an owner with the authority to act.

12.7 Production Debugging and Incident Response

At 3:00 AM, the paging system alerts the on-call engineer that revenue from the core recommendation system has dropped 15 percent in the last hour. The servers are healthy, the latency is normal, and there are no exception logs. Conventional software can also return incorrect results silently, but machine learning adds failures rooted in data and statistical behavior while ordinary service-health signals remain normal. Debugging production ML systems therefore requires investigative frameworks that reach beyond code and infrastructure.

At platform scale, production debugging can consume a large share of engineering attention, and its complexity multiplies because failures may originate in data pipelines, model code, infrastructure, or emergent interactions between components that no single team owns end to end. Effective incident response therefore classifies the symptom and contains impact while attribution proceeds, then converts the incident into a stronger platform control. A known safe containment action need not wait for complete root-cause analysis.

War Story 12.1: The control plane that disappeared (2021)

Context: Meta's global backbone connected its data centers and carried dependencies whose loss also made DNS and internal recovery tools unavailable (Janardhan 2021).

Mechanism: On October 4, 2021, a maintenance command intended to assess backbone capacity disconnected the data centers. A bug in the command-audit tool failed to stop it. DNS servers then withdrew their Border Gateway Protocol (BGP) advertisements, the routing announcements that tell other networks how to reach them, after losing data-center connectivity. DNS and network unavailability also disabled internal recovery tools and remote access.

Impact: The outage interrupted services worldwide and disabled internal tools needed to diagnose and resolve it, extending the consequences of the backbone failure into the recovery process itself.

Fix: Engineers went onsite to the data centers, completed secure access procedures, and restored backbone connectivity. They then brought services back under controlled load to avoid a second failure during recovery.

Systems lesson: Platform-scale operations must treat the control plane as a hard dependency with a defined blast radius. ML platforms present identical failure risks at fleet scale: training schedulers, model registries, feature stores, and inference routers share the underlying auth, DNS, and network planes. When network partitions occur, recovery requires bounded failure domains and an administrative path that remains usable when the normal control plane is unavailable. Physical separation is one possible mechanism, not a universal requirement.

12.7.1 Incident classification

The first five minutes of an outage are wasted if the responder restarts GPU serving containers for what is actually a stalled feature pipeline or a seasonal drop in user traffic. Incident classification is the first routing decision after an alert fires: the category tells responders which telemetry to inspect first, which owner to page, and whether the likely control is a rollback, a data repair, or infrastructure isolation. The incident categories are organized around that first diagnostic signal.

Data incidents involve problems with input data: pipeline failures that prevent fresh data from reaching models, schema changes that break downstream consumers, data quality degradation from missing values or distribution shifts, and feature staleness exceeding SLO thresholds. These incidents often manifest as accuracy degradation across multiple models that share data sources, making data pipeline health the first diagnostic checkpoint.

Model incidents involve problems with model behavior, including accuracy degradation beyond acceptable thresholds, latency spikes indicating computational issues, memory exhaustion from growing state (KV cache, buffers), and prediction bias shifts detected by fairness monitoring. Model incidents typically affect individual models. When multiple unrelated models degrade simultaneously, a shared data or infrastructure issue becomes the leading hypothesis rather than independent model problems.

Infrastructure incidents involve problems with the serving platform: GPU failures causing request errors, network partitions between model shards, load balancer misconfigurations routing traffic poorly, and container orchestration issues affecting deployments. These incidents tend to produce error rate spikes and timeout patterns rather than gradual accuracy degradation.

Business metric incidents involve unexpected changes to downstream KPIs, such as engagement drops without clear model or data cause, revenue anomalies during normal model operation, and user behavior shifts that affect model efficacy. These incidents are the hardest to attribute because they may stem from external factors (competition, seasonality, marketing campaigns) rather than ML system problems.

12.7.2 Attribution analysis

Once an outage is classified, attribution analysis tests candidate causes in sequence while containment limits user impact. When recommendation engagement drops 5 percent in an hour with no recent model deployments, rolling back the ranker wastes time if a 95 percent-full disk on a Kafka broker, one of the servers that store the event log feeding feature ingestion, is stalling ingestion three hours behind its freshness SLO. Responders can pause promotion, reduce exposure, or invoke a validated rollback before the root cause is certain, while preserving evidence for diagnosis. Temporal correlation traces the degradation backward through recent changes. In the trace, consumer lag counts the logged events that the feature pipeline has not yet read:

```
Symptom: Recommendation engagement dropped 5% in past hour

Step 1: Check recent deployments
        → No model deployments in past 4 hours
        → Lower the priority of a recent deployment; still check configuration and routing changes

Step 2: Check feature freshness SLOs
        → user_features: 3 hours stale (SLO: 1 hour)
        → Feature pipeline delayed

Step 3: Check feature pipeline status
        → Kafka consumer lag: 10M events (normal: 10K)
        → Data ingestion bottleneck

Step 4: Investigate Kafka cluster
        → Broker disk 95% full on partition 7
        → Storage pressure is a leading hypothesis, not yet a confirmed cause

Step 5: Test the storage-pressure hypothesis
        → Write stalls on the affected broker coincide with stalled ingestion
        → Bounded capacity recovery preserves committed data and consumer offsets
        → Consumer progress and feature freshness recover without a model change
```

The recovery evidence supports storage pressure as the ingestion bottleneck in this scenario. Responders still monitor engagement and delayed quality outcomes after freshness recovers; restoring pipeline progress does not by itself prove that all user impact has ended.

When measured accuracy drops, responders first verify the evaluation signal and then investigate changes in data, model behavior, and execution. Four recurring hypotheses within data and model incidents guide that investigation. Data drift is an input-distribution change, such as evolving demographics or seasonality, whose quality effect requires outcome evidence. Feature staleness arises when ingestion or pipeline delays leave predictions using outdated inputs. Model decay can reflect concept drift, where changed input-output relationships reduce the usefulness of learned parameters. Upstream model changes can invalidate representation assumptions across downstream

consumers. These hypotheses are not exhaustive; label defects, configuration changes, and numerical faults can also alter measured accuracy.

The diagnostic sequence then tests the highest-probability shared causes before local model defects:

1. Compare current input distribution to training distribution
2. Check feature freshness across all input features
3. Examine performance on stable evaluation sets
4. Trace dependency graph for recent changes

Example 12.3: Feature pipeline cascade

Scenario: A variant of the `user_engagement_score` change from section 12.4.4 ships without versioning. A core platform team alters the shared feature's transformation logic in place, changing its normalization from a 30-day z-score to a 7-day min-max scale.

Diagnosis: Fourteen downstream models across five engineering teams consumed the feature as an un-versioned interface. The in-place mutation invalidated feature distribution assumptions, causing a 12 percent drop in ranking accuracy and 18 hours of cross-team debugging toil.

Systems lesson: Shared feature definitions are production API contracts that require explicit, immutable versioning (`user_engagement_score_v2`) for incompatible changes. A transformation can preserve the field name and type while changing its meaning, so schema validation alone cannot protect downstream model assumptions.

At platform scale, the affected model set helps prioritize shared dependencies for investigation. Table 12.22 maps observed patterns to candidate causes; dependency records and component evidence must establish whether the suggested cause explains the incident. A single affected model can still depend on faulty shared infrastructure, while a broad pattern can reflect a common traffic change:

Table 12.22: Cross-Model Failure Patterns: At platform scale, the pattern of which models degrade simultaneously prioritizes shared dependencies to investigate. Correlated degradation across all recommender systems can motivate checking shared features; an all-vision degradation can motivate checking shared image preprocessing; a single-model degradation starts the investigation at that model and its dependencies.

Pattern	Likely cause
All RecSys models degraded	Feature store issue
All vision models degraded	Image preprocessing pipeline
Single model degraded	Model-specific issue
Geographic pattern	Regional infrastructure
Time-based pattern	Batch job scheduling

12.7.3 Runbook development

Because cross-model attribution requires checking shared dependencies before local model containers, a runbook that says "Investigate the issue" or "If BERT model fails, restart container" wastes the first half hour on healthy serving pods while a stalled feature pipeline keeps serving stale inputs. A runbook earns its place when it preserves diagnostic order under time pressure. If latency and error rates are nominal but business KPIs fall, responders inspect feature freshness, input distributions, and recent model or data changes before restarting infrastructure; if requests time out or a region becomes unreachable, serving and infrastructure health take priority. In either path, the runbook names a safe containment action, the evidence to retain, and the owner authorized to escalate.

The reusable structure is a diagnostic flow, not a document template. For a recommendation engagement drop, the responder should move through the same control loop every time: detect

user impact, localize the shared dependency, bound the blast radius, escalate by evidence, and learn from the gap (table 12.23).

Table 12.23: Runbook Diagnostic Flow: Symptoms prioritize dependency checks, and safe containment proceeds whenever the evidence warrants it. Attribution and mitigation can proceed concurrently; the runbook records evidence, authority, and the rollback path.

Step	Question	Example evidence
Detect user impact	Which product or cohort is degraded?	CTR, conversion, retention, revenue, or complaint rate
Localize dependency	Did a shared service change for many models?	Feature freshness, model registry, regional health
Bound blast radius	Can traffic be reduced, routed to a fallback, or rolled back?	Canary metrics, rollback target, serving saturation
Escalate by evidence	Which team owns the failed control?	Data pipeline, platform, serving, or model owner
Learn from the gap	Which missing signal would have caught it sooner?	Freshness gate, slice alert, deployment guardrail

Runbook anti-patterns break diagnostic order in three common ways: encoding a single historical fix rather than a reusable investigation, transferring all judgment back to the responder with vague instructions at the worst possible time, or pointing responders toward deprecated systems, dead dashboards, and obsolete contacts. Avoiding these failure modes requires treating the runbook as a maintained operational artifact that is exercised and updated alongside the service itself.

12.7.4 Post-incident reviews

Once the runbook restores service, the operational question shifts from immediate containment to preventing recurrence. Closing a feature-outage report with "the Kafka disk filled up" or blaming the engineer on call guarantees that the next stalled pipeline will reach production undetected, because no freshness gate blocked the stale features before user engagement dropped. Post-incident reports (PIRs) transform incidents into organizational learning only when they turn one failure into a stronger platform control. The incident record connects duration and user impact to the timeline of detection, attribution, mitigation, and recovery. Root causes identify the control that failed or did not exist, and corrective actions name owners and deadlines so the review changes the platform rather than only documenting the outage.

The useful PIR fields are the ones that force a control decision. Duration and impact quantify severity; the timeline exposes detection and attribution latency; root causes name missing controls; and corrective actions assign ownership for the next platform change. The questions in table 12.24 turn that timeline into platform work.

Table 12.24: PIR Control Questions: A post-incident review is valuable when it converts an incident timeline into a stronger platform invariant.

PIR question	Control decision it should force
How long was user impact visible?	Whether detection latency or mitigation latency dominated the incident.
Which dependency failed first?	Whether the missing control belongs to data, platform, serving, or model code.
Which signal arrived too late?	Whether a new freshness, slice, saturation, or canary guard is needed.
What would have bounded impact?	Whether rollback, exposure limits, admission control, or fallback was absent.
Who owns the next control?	Whether the review changes the platform rather than only recording blame.

Effective PIRs require psychological safety and a systems frame rather than individual blame. The review should identify which systems allowed the incident, not which person caused it. It should ask which signal would have detected the failure earlier, not why an individual missed it. It should prevent the entire class of failure, rather than merely the exact failure that occurred.

12.7.5 Incident ownership for distributed failures

Assigning corrective actions after an outage works only when runtime, infrastructure, and modeling teams also have clear authority during the outage itself. Otherwise, when a 1,024-GPU training job stalls at a collective barrier or diverges on step 4,000, the logs cannot be joined without shared rank, host, switch, and data-shard identifiers. Those teams then debate the cause, and no one has authority to quarantine a node or halt the run. Distributed training and inference can fail at a rank, device, host, fabric link, data shard, or control-plane dependency, and the fleet telemetry of section 12.6.1 collects the raw signals that detect such failures.

Diagnosing and recovering an individual fault belongs to the owning subsystem and its fault-tolerance mechanisms. Incident response adds a narrower contract on top of the fleet telemetry. It joins the signals by stable identifiers so responders can identify the affected domain, assigns authority before an incident occurs, and preserves the evidence needed for escalation and postincident learning.

↩ Heartbeat failure detection and in-flight job recovery procedures are established in section 7.9.

The telemetry contract starts with stable identifiers. Every event must carry the job or service identifier, attempt and model version, rank and host, device identity, topology location, data-shard or request cohort, software build, a synchronized wall-clock or trace timestamp for cross-host correlation, and a host-local monotonic time or sequence number for local duration and ordering. Metadata naming which clock produced each timestamp, and how closely that clock is synchronized with other hosts, bounds how precisely events from different hosts can be ordered. Logs without these dimensions force responders to reconstruct joins during an outage; unbounded high-cardinality labels can overwhelm the monitoring system. The platform therefore stores low-cardinality health aggregates for alerting and retains sampled or event-scoped high-cardinality traces for diagnosis. Correlation, not raw log volume, determines whether telemetry is operationally useful. With events joined on these identifiers, each class of failure evidence has a first owner and a defined escalation record (table 12.25).

Table 12.25: Distributed Failure Telemetry and Ownership: The platform collects correlated evidence and routes the incident; the owning subsystem team performs mechanism-level diagnosis. Stable identity and topology fields let responders distinguish one defective rank from a fleet-wide software or fabric failure.

Failure evidence	Fleet-level signal	First owner	Escalation evidence
Collective progress stalls	Ranks entering and completing each collective by topology domain	Training runtime	Blocked rank set, last completed collective, fabric locality
One rank approaches its memory limit	Per-rank allocated, reserved, and peak memory distribution	Model/training team	Allocation phase, model state, comparison with peer ranks
Step time develops a long tail	Rank-level step-time distribution correlated with host and switch	Fleet infrastructure	Straggler persistence, thermal, network, and input evidence
Loss or gradients diverge across ranks	Rank-specific numerical summaries tied to artifact and data shard	Model/training team	First divergent step, shard, build, and checkpoint
Repeated attempts fail on one domain	Failure rate by host, rack, device generation, and software build	Fleet infrastructure	Quarantine history and reproduction on a known-good domain

Ownership must include the authority to contain impact. The training-runtime owner may cancel or restart a job when progress stalls; fleet infrastructure may quarantine a device, host, or topology domain; the data owner may stop a corrupted shard; and the model owner decides whether numerical divergence invalidates the resulting artifact. The incident commander coordinates these actions when more than one domain is implicated. A dashboard that identifies a likely cause without naming who may stop traffic, quarantine capacity, or invalidate an artifact leaves the most consequential decision unresolved.

Telemetry retention closes the control loop. Routine healthy runs need aggregate distributions and sampled traces, whereas failed attempts and release-gate violations require full event-scoped evidence through the postincident review and corrective-action window. The platform should measure attribution latency, percentage of incidents routed to the correct owner on first escalation,

recurrence by failure class, and time from corrective action to deployed control. Those measures reveal whether observability is reducing fleet risk or merely generating more data.

12.7.6 On-call practices for ML teams

Even clear ownership tables and correlated identifiers fail if the human rotation carrying the pager is structured for traditional web uptime alone. Such a rotation asks one responder to diagnose a GPU scheduler, trace a stale feature pipeline, and interpret a statistical drop in prediction quality during a single shift. ML systems therefore require specialized on-call practices because prediction quality, data freshness, and dependency graphs join uptime as reliability concerns. The attribution analysis of section 12.7.2 shows why an incident can begin as a business metric anomaly and end in a feature pipeline, a model registry, or a stalled distributed operation. On-call design has to make that diagnostic burden sustainable, building on established site reliability engineering (SRE)[13] principles (Beyer et al. 2016).

Because data and statistical failures expand the evidence responders must interpret across platform, data, and modeling domains, a tiered response assigns initial containment, triage, and escalation to a primary responder, with clear paths to domain specialists. A 15-to-30-minute routing target is an illustrative operational objective to calibrate against incident severity and coverage, not a universal deadline or a reason to delay an immediately available containment action.

Rotation design determines whether the team can sustain that diagnostic burden. A reasonable starting policy is a one-week primary rotation with a secondary backup and a 30-minute handoff overlap, adapted to page volume, staffing, time zones, and incident duration (table 12.26). Secondary coverage supplies backup expertise, while explicit handoffs preserve context. Rotation length alone cannot prevent burnout if interruptions and recovery work exceed the team's capacity.

13 **SRE**: Founded by Ben Treynor Sloss at Google in 2003, SRE introduced error budgets and SLOs as quantitative reliability contracts. For ML systems, classical SRE must be extended. Traditional services can fail semantically as well as become unavailable, and models add statistical quality degradation that an uptime check cannot detect. This demands drift-aware SLOs that treat prediction quality alongside uptime as a reliability metric.

Table 12.26: On-Call Rotation Design: An illustrative starting policy for rotation length, backup coverage, and handoffs. Interruption burden, recovery work, staffing, and context-transfer quality determine whether the policy is sustainable; rotation length alone does not establish burnout risk.

Aspect	Recommendation
Rotation length	Start with 1 week; adapt to page load, staffing, and recovery work
Primary + secondary	Always have backup; ML incidents often require multiple experts
Handoff overlap	30 min overlap for incident context transfer
Follow-the-sun	Target 8-hour regional shifts, with explicit handoffs and backup

Alert fatigue, the cumulative false-alarm burden of section 12.6.2, shows up on call as discounted notifications, slower acknowledgment, and pages with no useful response. Teams therefore recalibrate thresholds against observed false alarms and missed incidents on a declared schedule, track the share of pages that lead to action, and require each page to name its owner, diagnostic evidence, and runbook.

Beyond general SRE skills, ML on-call requires interpreting model quality metrics, understanding data dependencies, and choosing safe containment while diagnosis continues. Toil reduction matters because recurring manual tasks consume the capacity needed to improve the platform, so teams track their share of on-call time over a declared measurement window. A target below 25 percent is one planning choice, to be assessed alongside page load, recovery work, and time available for engineering improvements.

Common ML toil accumulates at the model-platform boundary in tasks such as restarting failed runs, processing routine deployments, triaging repeated alert noise, and compiling reports. Automation can reduce this overhead when the action and its failure conditions are understood, and its payback follows the same calculation as notebook 12.4, extended by validation, maintenance, exception handling, and the risk of automating the wrong response.

The controls built so far, from release pipelines, registries, and feature contracts to hierarchical monitoring and incident response, are expensive to build, and building them separately for every product team would recreate the duplication that motivated them. Organizations therefore package these capabilities as a unified internal product, a discipline known as platform engineering.

12.8 Platform Engineering

Retiring the recurring toil surfaced by on-call rotations requires turning ad-hoc operational fixes into shared infrastructure. Organizations where every data science team independently provisions GPU clusters, configures model registries, and wires up alerting dashboards duplicate that infrastructure work across the fleet. **Platform engineering** addresses this by treating the ML infrastructure itself as an internal product, providing supported paths that absorb recurring operational complexity while exposing the controls model teams need. Product teams still own model behavior and workload-specific decisions, while the platform balances self-service capabilities that accelerate development against governance requirements that enforce common release and reliability policies.

12.8.1 Abstraction levels

Handing ten model teams raw Kubernetes namespaces forces each team to rebuild its own training loop, serving stack, and monitoring dashboard, whereas locking every team into a rigid end-to-end wizard blocks the first workload that needs a custom distributed topology. The abstraction-level decision is how much operational judgment the platform should centralize on behalf of model teams. The four levels in table 12.27 sit on that flexibility-vs.-convenience curve, distinguished by which concerns the platform owns and which it leaves to the model team.

Table 12.27: Platform Abstraction Levels: Platform engineering centralizes different parts of the ML lifecycle at each abstraction level. The right level depends on whether an organization values workload-specific control or shared operational invariants more.

Level	Platform owns	Model team still owns	Best fit and risk
Level 1	GPU capacity, storage volumes, network connectivity, and basic orchestration such as Kubernetes namespaces	Training code, serving stack, monitoring, and deployment path	Fits unusual workloads with strong infrastructure teams; becomes costly when many teams repeat the same operational work
Level 2	Standardized containers, ML-aware Kubernetes integration, persistent volumes, and basic service-mesh support	Experiment structure, training jobs, serving releases, and monitoring workflows	Reduces packaging duplication; operational correctness still depends heavily on each team
Level 3	ML-specific control loops such as topology-aware scheduling, hyperparameter tuning, distributed training, serving	Modeling judgment and workload-specific trade-offs	Absorbs failure modes that recur across teams; systems such as Kubeflow and Ray often sit near this level
Level 4	Full lifecycle product: IDEs, feature stores, experiment tracking, registries, CI/CD, monitoring, cost, governance	Higher-level product and modeling intent, with less direct low-level control	Maximizes speed and consistency; the platform team accepts responsibility for policy, reliability, and economic visibility

Managed services such as Vertex AI and SageMaker, and lifecycle tooling such as TensorFlow Extended (TFX)[14] at Google (Baylor et al. 2017) and MLflow (Zaharia et al. 2018), support different parts of this spectrum. A product name does not establish the organization's abstraction level. The decisive boundary is which lifecycle responsibilities the assembled platform actually owns, which controls model teams retain, and who supplies policy, reliability, and economic visibility for the gaps.

[14] **TensorFlow Extended (TFX)**: Google's production ML pipeline framework includes TensorFlow Data Validation (TFDV), which checks data against schema expectations and configured drift and skew comparisons. Connecting those checks to a release policy allows a pipeline to stop before spending training resources on detected data defects. The checks cover declared expectations; a passing result does not establish that every value is correct or that a distribution change is harmless.

12.8.2 Platform lifecycle and adoption

When a platform team treats its internal stack as a one-time infrastructure build, every breaking upgrade forces product teams into a synchronized migration that stalls launches, or leaves behind "shelfware" with high signup counts where teams still maintain custom shadow pipelines. An internal ML platform is a versioned product whose interfaces, supported workflows, reliability targets, and ownership boundaries must evolve without halting model delivery. Each platform release therefore needs compatibility guarantees, a migration path, telemetry that distinguishes old and new paths, and a retirement date for deprecated behavior. The platform team also needs its own service-level objectives because a registry, feature service, or deployment controller outage can block every product team at once.

Adoption measures whether the platform removes work rather than merely centralizing it. Useful indicators include the share of production models on supported paths, time from validated artifact to bounded production traffic, exception rate, migration completion, support load, and repeated toil outside the paved road. A high signup count with persistent custom pipelines signals nominal adoption but weak product fit. Conversely, mandatory migration can raise the adoption percentage while increasing lead time and operational risk. The platform roadmap should therefore combine fleet outcomes with user evidence: incident reduction and cost attribution matter alongside time-to-first-deployment and the fraction of teams that can complete routine changes without platform intervention.

Lifecycle decisions must preserve an escape hatch for workloads whose constraints do not fit the shared path. Exceptions should be explicit, time-bounded contracts with named owners and integration obligations, not permanent forks. When exception volume grows around the same missing capability, the pattern becomes product evidence for the next platform version. When a capability no longer earns its maintenance and migration cost, the platform should deprecate it through measured dual-running, compatibility testing, and a final removal gate. This product discipline keeps the shared control plane from becoming the next source of technical debt.

12.8.3 Self-service model deployment

If every model release requires a ticket to the infrastructure team, deployments stall in a central queue; if model teams push directly to production without guardrails, a single unvalidated config can exhaust cluster GPU quotas or bypass canary gates. Self-service deployment works only when the platform lets model teams move quickly while keeping risky choices inside governed boundaries. The design question is therefore which deployment decisions remain team-owned and which invariants the platform enforces before a model receives production traffic.

A well-designed deployment API abstracts operational complexity by making the model team's intent declarative. The platform lesson is not the shape of a particular YAML file; it is which invariants the platform must capture before it turns a model version into fleet traffic. The invariants in table 12.28 are the platform contract behind that declarative interface. Prometheus[15] is one common operational sink for the telemetry side of that contract.

Table 12.28: Self-Service Deployment Invariants: A deployment API should encode the controls that keep local team autonomy from creating fleet-wide risk. The concrete syntax can vary, but the invariants must survive across platforms.

Declared intent	Platform invariant
Model artifact and version	The serving system loads exactly the validated model from the registry.
Resource envelope	GPU, memory, and replica requests stay inside quota and scheduling constraints.
Traffic policy	Canary, shadow, and rollback controls bound blast radius before full promotion.
Quality gates	Error, latency, drift, and task-quality thresholds block unsafe promotion.
Telemetry sink	Operational metrics flow into operational monitoring, while model-quality signals flow into drift and slice monitors.
Approval boundary	Sensitive data access, budget overruns, and high-risk release windows require review.

TensorFlow Serving provides model-serving lifecycle capabilities for loading, versioning, canarying, and rolling back models (Olston et al. 2017). A self-service platform adds the policy layer around those capabilities: model teams specify intent, and the platform translates it into scheduling, traffic routing, telemetry, and approval controls.

12.8.4 Resource management

Packing batch training jobs onto a shared cluster to hit 95 percent GPU utilization leaves co-located inference endpoints without headroom when user traffic spikes, while preempting a 72-hour training job that lacks a recent checkpoint wipes out days of compute. Training jobs can trade completion time for utilization within their deadlines; serving needs headroom to protect tail latency as requests arrive. A shared scheduler can manage both, provided its policy represents those distinct constraints rather than maximizing utilization alone.

[15] **Prometheus**: An open-source monitoring system that scrapes and stores labeled time series. A fleet can export operational metrics and computed drift scores through the same metrics interface while retaining detailed prediction records elsewhere for analysis. The distinction is between compact alerting signals and the evidence needed to investigate them; it does not require every platform to adopt the same two-tier implementation.

[16] **Job preemption**: The ability to terminate lower-priority workloads to free resources for urgent work. Preemptible capacity is also how cloud providers sell deep discounts, which section 12.9.4 prices. Without a usable checkpoint, a preempted 72-hour run may need to restart from scratch, potentially erasing the discount. Checkpoint frequency trades I/O overhead against wasted-compute risk; the right policy depends on job duration, interruption frequency, and recovery cost.

Training resource policy must expose which jobs may wait or be preempted,[16] which deadlines are binding, and who accepts the recovery cost. Scheduler mechanisms enforce those choices. The platform records checkpoint readiness, budget, priority, and accountable owner with each job so a utilization improvement does not silently transfer deadline or lost-work risk to a model team.

↩ Scheduler mechanisms for queueing, priority preemption, and elastic resizing of training jobs are established in chapter 8.

Serving binds on latency rather than batch flexibility, because a live request cannot be preempted without user impact. Its control loop therefore starts from the latency SLO and the capacity that must be online before requests arrive. Autoscalers translate request load and queue pressure into capacity decisions; operations owns their policy envelope: the latency SLO, warm rollback capacity, tenant limits, budget, and authority to override normal scaling during an incident. Those commitments must remain visible in the release record so a cheaper allocation cannot silently remove the headroom on which validation relied.

↩ Autoscaling controllers that convert queue growth and token arrival rates into replica counts are established in section 10.13.

12.8.4.1 Platform utilization metrics

Once training and serving controls are separated, the shared platform still needs an aggregate measure of whether capacity is being converted into useful work. Equation 12.11 defines platform efficiency as the capacity-weighted average utilization across all resources:

$$U_{\text{platform}} = \frac{\sum_i U_i \times \text{Cap}_i}{\sum_i \text{Cap}_i} \tag{12.11}$$

where U_i is the utilization of resource i and Cap_i is its capacity in a common unit over the same interval. Heterogeneous accelerators require a declared workload normalization; CPU cores, device counts, and bytes cannot be added directly. Weighting by capacity keeps the metric honest about where stranded work lives. An idle single GPU and an idle eight-GPU node are not equally wasteful, so a small underused resource should not drag the platform average down as hard as a large idle pool.

However, raw utilization is incomplete. The scheduler already separates allocated, compute, and productive utilization for a single cluster; the platform metric needs the same distinction. Effective utilization must also weigh utilization quality, whether busy GPUs are doing productive work or waiting on data; utilization fairness, whether each team receives capacity in line with its agreed allocation; and utilization cost, the cost per unit of ML output.

↩ The allocated, compute, and productive utilization tiers are established in section 8.12.1.

12.8.4.2 GPU cluster efficiency

A platform operates a 100-GPU training cluster with average GPU utilization of 65 percent, memory utilization of 80 percent, four-hour average queue waits, and an internal chargeback rate of \$2.50/GPU-hour, the price the platform bills teams for each GPU-hour. Memory occupancy and compute utilization alone do not identify the bottleneck. Profiling must separate input stalls, communication, kernel inefficiency, and scheduling gaps, and queue waits may reflect burst arrivals or placement constraints rather than insufficient aggregate capacity. Busy time is also not useful work. The cluster costs $100 \times 24 \times \$2.50 = \$6,000$/day at the chargeback rate, and its busy share, $100 \times 24 \times 0.65 \times \$2.50 = \$3,900$/day, is only an upper bound on the value of its useful output. Busy accelerators can still run jobs that later fail, repeat work lost after a checkpoint, or stall inside a step. The goodput accounting of section 9.2.3 charges each of those losses to the layer that caused it. Suppose profiling attributes the gap to input stalls and a data-loading repair raises busy time to 80 percent, $100 \times 24 \times 0.80 \times \$2.50 = \$4,800$/day, with throughput validated at the same quality target. The fixed cluster bill does not change; whether the recovered capacity becomes cash, by avoiding a purchase, or value, by serving queued work, follows the utilization valuation of section 8.8. Scheduling and capacity planning then address the remaining queue pressure.

> **Checkpoint 12.4: The annual cost of GPU cluster underutilization**
>
> Use daily cluster spend (\$6,000/day) and busy-capacity value, utilization times spend and an upper bound on useful work (\$3,900/day at 65 percent utilization), to extend the analysis to annual scale and compare against the platform ROI argument.

- ☐ **Underutilization cost**: calculate the annual cost of underutilization for the 100-GPU cluster at 65 percent utilization, the cost after optimization to 80 percent, and the annual value of recovered capacity. State when that value would become cash savings.
- ☐ **Queue dynamics**: Little's law (theorem 10.1) holds for any stable queue, rather than serving requests alone. Applied to the training-job queue, the average number of jobs waiting equals the job arrival rate times the average queue wait. Calculate the job arrival rate for a 4-hour average wait with 100 jobs in queue and explain what additional job-size, service-time, and arrival-pattern measurements are needed to choose expansion or utilization optimization.
- ☐ **Savings vs. platform**: compare the annual recovered-capacity value from cluster optimization to the $2M/year platform investment scenario. Holding the utilization gain and chargeback rate fixed, determine the GPU count where annual recovered-capacity value equals that investment, and state what must happen for the capacity value to become cash savings.

12.8.5 ML productivity goodput across the portfolio

A GPU spinning at 100 percent utilization on a hyperparameter tuning job that eventually fails from a configuration error looks efficient on a hardware dashboard while producing zero engineering value. Raw utilization metrics capture device busyness, not useful progress; ML productivity goodput (MPG) turns that distinction into a fleet metric (Wongpanich et al. 2025).

MPG multiplies scheduling, runtime, and program goodput, which measure how much capacity survives each stage of a job's life: holding a complete allocation, retaining progress while it holds one, and running its retained steps at the ideal rate. Each lost hour is charged to exactly one factor, so the product says which layer limits useful output as well as how much output was lost. The platform's job is to compute the three factors per team and workload class from scheduler records, job timelines, and step traces. A portfolio dashboard then shows which layer limits each team and which owner, orchestration, the runtime, or the program, must act.

↩ ML productivity goodput, its three factors, and the boundary at which each one counts a loss are established in section 9.2.

12.8.6 Multi-tenancy and isolation

When a batch training job claims all free HBM on a shared node and starves co-located production inference, or when a shared cluster generates a monthly GPU bill that no team can attribute to its own workloads, multi-tenancy fails before its efficiency savings can materialize. The scheduler's quotas and fair-share policies and the serving system's isolation mechanisms, established in section 10.12, already bound interference within one cluster or one inference service. What the platform adds is one namespace hierarchy across training and serving, plus cost isolation that follows each tenant's spend, and neither the scheduler nor the serving system owns those alone. Multi-tenancy is therefore an isolation problem before it is an efficiency problem, and the platform must bound performance interference, data exposure, and cost spillover at the same time.

The platform's added obligations are hierarchy and attribution. A typical design gives each team its own development, staging, and production namespaces with resource quotas, and places shared services such as the feature store, model registry, and monitoring in common namespaces with access controls, network segmentation, and encryption at the tenant boundary. Cost isolation then combines metering with declared allocation rules for shared services and reserved capacity, because not every shared cost can be assigned to an individual request. Without these boundaries, one team's demanding workload can degrade performance for the others, the noisy neighbor problem, which on an ML platform crosses the boundary between training and serving whenever an unthrottled training job contends with co-located inference. The request limits, rate limits, priority classes, and burst budgets that enforce fairness in the data plane are the serving and scheduling mechanisms already cited; the platform's job is to apply them consistently across every tenant's namespaces.

↩ Quotas, hierarchical fair-share, and burst borrowing are established in section 8.11.

Isolation makes sharing safe, but it does not make sharing pay. Cost isolation meters each tenant's consumption; whether the shared fleet is worth its capital, how its spend reaches the teams that caused it, and when a workload's cost justifies optimization are economic questions that the platform must answer with the same rigor it applies to releases.

↩ Noisy-neighbor interference and the device-sharing mechanisms that bound it are established in section 8.10.1.

12.9 Fleet Economics and Utilization

A provisioned GPU incurs capacity costs whether it serves predictions, waits between training bursts, or remains reserved for rollback. Those idle periods have different meanings: a scheduling bottleneck strands capacity, while a justified reserve buys recovery capability. The operational machinery developed so far decides how models change safely; it does not decide whether the capacity underneath them is worth what it costs. Fleet economics therefore distinguishes productive use, deliberate headroom, and avoidable waste before treating higher utilization as an improvement.

Three economic levels keep the cost analysis from collapsing unlike decisions into one number. *Platform economics* (section 12.2.3) asks whether common control-plane capabilities repay their fixed development and maintenance cost across the portfolio. *Fleet economics* asks whether shared hardware, facilities, and capacity commitments produce useful work at acceptable utilization and risk. *Workload economics* asks whether one training run, serving endpoint, or experiment earns its marginal compute, data, and engineering cost. A gain at one level does not establish value at the others. Faster jobs do not remove an existing capacity commitment, and high fleet utilization does not establish that a platform repays its adoption and maintenance costs. The placement and scheduling mechanisms in chapter 8 operate within this capacity envelope. Operations owns the commitments around them: capital planning, portfolio return, lifecycle total cost of ownership (TCO), and cost attribution across teams.

Definition 12.1: ML systems TCO

Total cost of ownership (TCO) for ML systems is the complete economic accounting of developing, deploying, and operating machine learning capabilities across their full lifecycle.

1. **Significance**: It captures the cost inversion of scale: training incurs a cost per run, while inference accumulates with served work. The lifecycle example in section 12.9.2.2 prices that inversion for a high-traffic service, where cheaper serving repays additional training spend.
2. **Distinction**: This lifecycle view includes serving efficiency alongside hardware, data, development, and recurring training costs. The energy and compute cost of each prediction depends on achieved hardware efficiency η_{hw} as well as batching and model size.
3. **Common pitfall**: A frequent misconception is that the initial training bill is the primary financial risk. In reality, the data debt and maintenance overhead are the "silent interest rates" that can make an accurate model economically unsustainable in production.

Platform economics was priced earlier in the chapter, so the analysis here follows the remaining levels downward. It prices the fleet first, then the workload, and closes with financial operations (FinOps), the controls that attribute those costs to teams and assign the authority to act on them.

Engineering constraints shape the feasible economic choices. Infrastructure must first meet quality, latency, reliability, and capacity requirements; useful computation per dollar then helps compare feasible designs over the system's lifetime. At an assumed $350,000 per 8-GPU DGX H100 node, a 10,000-GPU cluster represents about $437.5M in node hardware capital expenditure before networking, facility, staffing, and maintenance costs. The purchase price, however, is only the beginning. Over the three-year planning horizon used here, those operating and supporting costs determine the full commitment. TCO[17] connects that commitment to useful output and provides the basis for the broader ML-lifecycle accounting developed in section 12.9.6.

An accelerator fleet concentrates a large commitment in resources that stand idle whenever data, communication, or scheduling falls behind, so the economic target is useful output within the headroom that failures, rollbacks, and traffic spikes require, not maximum device busyness.

Two reference prices anchor the arithmetic that follows. The internal fleet rate of $2/GPU-hour approximates what an owned H100 fleet operated at high utilization costs its own teams, and the on-demand cloud rate of $4/GPU-hour is the rental reference for build-vs.-buy comparisons. Other rates in the chapter, such as the internal chargeback price in the GPU cluster efficiency example, the FinOps spending scenario, or the non-energy accelerator-hour rate in the training cost example, are scenario choices and are labeled as such where they occur. Cluster sizes follow the same rule. The

17 **TCO boundary**: A fleet comparison must declare what its total cost of ownership counts. The boundary must include the supporting power, cooling, networking, facility, and staffing commitments. Moving from 20 percent to 80 percent useful utilization spreads fixed costs over more output, but whether that changes the ownership decision depends on the rental alternative, workload fit, and costs included in each option.

running fleet is the Archetype A reference cluster of section 1.6.1.1, the H100 cluster that trains the GPT-3-class reference instance; every later analysis prices it unless it names a separate scenario.

12.9.1 Cost centers as operating constraints

Standing up the running 1,024-GPU H100 cluster (128 nodes) commits $44.80M in server hardware before a single job runs, plus $5 million for the InfiniBand fabric, $10 million for facility buildout, and $1.5 million each year in power and staffing. The total cost of an ML cluster decomposes into two broad categories. **Capital expenditure (CapEx)** covers the one-time costs of building the infrastructure: accelerators, servers, networking equipment, facility construction, and installation. **Operational expenditure (OpEx)** covers the recurring costs of running it: electricity, cooling, network bandwidth, staffing, maintenance, and software licenses. For a large on-premises cluster, table 12.29 separates the cost stack into capital and operating centers, each with different utilization implications.

Table 12.29: Frontier Training Cost Stack: Capital costs are dominated by accelerators, networking, and facilities, while operating costs are dominated by electricity and specialized staffing. The shares are broad planning ranges; in the running 1,024-GPU example, node hardware is about 75 percent of CapEx and electricity about 44 percent of OpEx.

Cost center	Cost class	Typical share	Included components and implications
Accelerators and servers	CapEx	50–60% of CapEx	GPUs or TPUs, host servers, baseboard management controllers, and local storage; this category dominates upfront purchase and depreciation.
Networking	CapEx	10–15% of CapEx	InfiniBand switches, HCAs, cables, and optical transceivers; the running example budgets about $5 million for a two-tier fat-tree serving 1,024 GPUs.
Facilities and cooling	CapEx	15–25% of CapEx	Building construction or retrofit, transformers, UPS systems, PDUs, and cooling plant; liquid cooling adds 10–15% to facility costs but reduces long-term OpEx.
Electricity	OpEx	60–70% of OpEx	At $0.07/kWh, a 1,024-GPU H100 cluster consuming about 1.1 MW including cooling at power usage effectiveness (PUE) 1.1 costs about $665,000 per year.
Staffing and maintenance	OpEx	20–30% of OpEx	System administrators, hardware technicians, replacement parts, and software license fees.

The running cost example is Archetype A's reference instance, the 175B-parameter GPT-3-class model, trained on the reference cluster of 1,024 H100 GPUs spread across 128 nodes. For the illustrative training plan developed below, assume a 2–4-week run. Parameter count alone does not determine this duration; the later planning example states its token and throughput budget.

Over a three-year lifecycle, whether owning or renting the cluster costs less depends on utilization. The hardware CapEx of $44.80M ($350,000 per node), the $5 million two-tier InfiniBand fat-tree network, the $10 million facility allocation, and $1.5 million per year in electricity (at $0.07/kWh with a PUE of 1.1) and specialized staffing bring the three-year total to roughly $64 million. If this dedicated cluster only trains six large models per year, the effective cost per run is $3.6 million. Renting the same 1,024-GPU allocation in the public cloud at $4/GPU-hour costs about $1.4M–2.8M per run (1,024 GPUs $\times$ 336–672 hours $\times$ $4), or $25M–50M over three years for six runs per year. Cloud rental is cheaper for this bursty scenario when idle allocations can be released. At the opposite extreme, *continuous utilization* spreads the owned cluster's stated costs over every available hour, reducing the cost to approximately $2.39/GPU-hour, below the cloud reference rate. Ownership need not require continuous use; the actual break-even lies where useful demand repays the fixed commitment under matched cost boundaries.

12.9.2 Utilization as the economic invariant

An organization that buys a dedicated cluster for six bursty training runs a year pays more per run than public cloud rental, whereas keeping that same cluster busy above its break-even utilization

↪ Power usage effectiveness as a facility metric, and what it omits, is developed in section 15.4.1.3.

turns upfront capital into a substantial per-GPU-hour savings. The most consequential economic question is whether the fleet can sustain enough useful load to amortize fixed capacity. Build-vs.-buy is one expression of that question, but the same invariant appears inside a deployed platform: reserved capacity, fallback pools, regional replicas, and canary headroom all become economical only when their reliability value justifies their idle cost. Ownership pays a large CapEx step up front and then climbs a shallow operating slope, while cloud rental starts at zero and rises steadily. Under sustained high utilization the rental line overtakes the owned cluster within a 30-month horizon (figure 12.13).

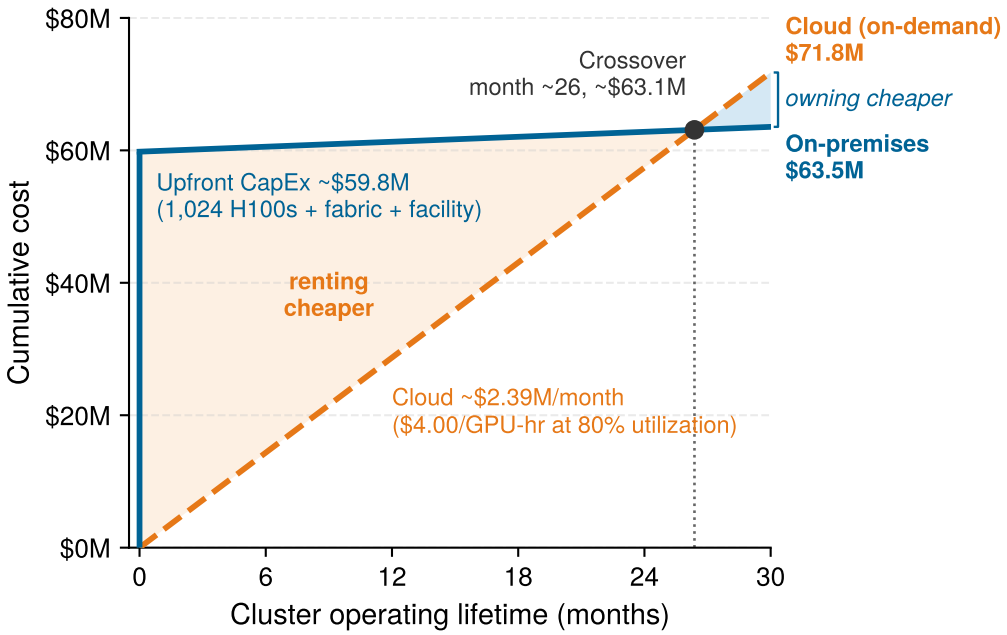

Figure 12.13: Build vs. Buy TCO: The curves compare cumulative cost over a 30-month horizon for a high-utilization H100 reference cluster. On-premises deployment begins with a CapEx step plus monthly OpEx, whereas cloud rental cost rises linearly with utilization. Over a 3-year service life, owning pays off only above 59.7 percent sustained utilization; below that, renting stays cheaper.

The plotted reference uses 1,024 H100 GPUs across 128 nodes at 80 percent sustained utilization. The on-premises curve starts at $59.80M in upfront CapEx and adds about $1.5M per year, while cloud rental at $4/GPU-hour grows by about $2.39M per month. Over 30 months, the cloud line reaches $71.8M, and the on-premises line crosses it around 26.4 months.

Cloud providers charge by the GPU-hour. At the on-demand reference rate of $4 per H100-hour, a single 8-GPU node rented only during 80 percent of the year costs $224,256 per year. An on-premises DGX H100 node costs approximately $350,000 to purchase. Amortized over three years and combined with electricity costs of $4,519 per node per year, the total annual on-premises cost is approximately $121,186 per node. This assumes idle rental periods can be released; an allocated but idle GPU still incurs a bill. Solving the break-even equation with electricity scaling at the same allocation fraction, on-premises infrastructure becomes favorable when sustained utilization exceeds roughly 42.5 percent. This narrow comparison includes node acquisition and utilization-proportional electricity; it excludes idle power, network and facility capital, staffing, maintenance, and financing. It does not use the wider owned-cost boundary of the cluster comparison in figure 12.13.

The same calculation explains why cloud pricing cannot be read as a simple price list. The scenario's reserved discount of 40–60 percent converts optionality into commitment. Its spot or preemptible discount of 60–90 percent introduces interruption risk. The relevant comparison is cost per completed workload under the deadline or service contract. Checkpoint recovery, replacement

capacity, and traffic fallback offer different ways to manage that risk; no single combination is required for every workload.

Owned capacity moves a different set of risks into the platform. Facility construction, networking, staffing, maintenance, and hardware obsolescence continue to accrue even when the GPUs are idle; their cost depends on site, power delivery, cooling design, and accounting boundary. A three-year-old GPU may be economically stranded if the next generation delivers 3× more performance per watt. The engineering question is therefore not whether one purchasing channel is universally cheaper; it is whether the workload mix can keep fixed capacity useful enough to compensate for its lifecycle and operating risks.

Hybrid designs follow from the same invariant. Owned infrastructure can carry predictable baseline work, while cloud allocations absorb peak demand, early hardware access, or bursty training campaigns. The hybrid split is not a compromise between two vendor categories; it is a scheduling policy over capital risk, utilization, and deadline risk.

For the 175B model, workload cadence changes the ownership decision. Consider a dedicated training fleet used only 15–20 percent of the time because its annual training campaign leaves long gaps. A portfolio of experiments, hyperparameter sweeps, fine-tuning jobs, and model variants could instead sustain 70–80 percent utilization on the same hardware, provided those jobs fit its capacity and scheduling constraints. Serving can also fill gaps when the hardware and isolation policy support it. These illustrative utilization levels describe workload mixes, not consequences of training a particular model. A hybrid design can keep predictable baseline work on owned capacity and rent the temporary expansion, such as the 5–10× baseline allocation assumed here for large training runs. The same invariant governs the comparison at a larger fleet scale, provided each estimate states which costs it includes.

Napkin Math 12.9: The 10,000-GPU cluster

Consider a cluster of 1,250 DGX H100 nodes (10,000 GPUs) for training the 175B model.

On-premises (3-year lifecycle):

- **Hardware CapEx**: 1,250 × $350,000 = $437.5M
- **Network CapEx**: ~$25M (InfiniBand fat-tree fabric)
- **Facility CapEx**: ~$75M (liquid-cooled data center hall)
- **Annual Electricity**: Simplified GPU-only estimate: 10,000 GPUs × 700 W × PUE 1.1 × 8,760 h/year × $0.07/kWh = $4.7M/year
- **Annual Staffing**: ~$5M/year
- **3-year total**: $537.5M + 3 years × $9.7M = ~$566.7M

Cloud rental (3-year rental at 80 percent utilization):

- **Annual Cost**: 10,000 GPUs × $4/GPU-hour × 8,760 h/year × 0.80 = $280.3M/year
- **3-year total**: ~$841.0M

Systems insight: With the stated prices and cost boundary, on-premises saves approximately $274.3M over three years at 80 percent utilization, but the savings disappear below roughly 53.9 percent. Owned hardware still incurs facility and staffing costs regardless of load. This is a scenario comparison, not a procurement quote: its GPU-only electricity estimate omits host and network power, and omitted maintenance or financing costs would also change the break-even point.

The three build-vs.-buy scenarios yield three different break-even utilizations because each draws a different cost boundary around owned capacity (table 12.30). All three use the same on-demand cloud reference rate, so the spread comes entirely from what each scenario charges to ownership, and the 10,000-GPU threshold shows how sustained demand, rather than fleet size, decides the ownership case.

Table 12.30: Build-vs.-Buy Break-Even by Cost Boundary: Three scenarios compare owned capacity with the same on-demand cloud rate. The single-node comparison omits network, facility, and staffing costs, so its threshold is lowest; the two cluster scenarios include them and differ by scale and by how each prices operating cost. A procurement decision needs the boundary of the cluster scenarios or wider.

Scenario	Scale	Owned cost included	Break-even utilization
Single-node comparison	One 8-GPU DGX H100 node	Node purchase amortized over three years; electricity proportional to use	42.5%
Build-vs.-buy figure (figure 12.13)	1,024 GPUs (the running cluster)	Nodes, network, and facility CapEx plus annual OpEx over the 3-year life	59.7%
10,000-GPU notebook	10,000 GPUs	Nodes, network, and facility CapEx; GPU-only electricity at full load; staffing; three years	53.9%

Systems Perspective 12.3: Sustained demand and unit cost

The 10,000-GPU scenario illustrates how sustained demand can support ownership. A portfolio that keeps compatible capacity usefully occupied at 70–90 percent utilization spreads fixed costs over more work than one operating at 20–30 percent. These are illustrative demand regimes, not utilization guarantees based on organizational size. Continuous training pipelines, model refreshes, and compatible inference workloads can support the higher regime; sporadic campaigns favor the option to release rented capacity. When a broad workload portfolio lowers unit cost enough to fund further investment, the advantage compounds, but it depends on useful demand and operating discipline, not fleet size alone.

12.9.2.1 Operational complexity

A TCO estimate must include the operational work that sustains the fleet. Hardware maintenance must diagnose failed GPUs, links, network adapters, and cooling components before a local fault becomes a fleet-wide slowdown. The five-hour fleet failure interval in the scenario of section 12.6.1 informs staffing, but does not determine headcount by itself. Staffing also depends on diagnosis and repair time, automation, concurrent incidents, and coverage requirements while the rest of the fleet keeps working.

The software side creates the same coupling. CUDA, GPU drivers, InfiniBand drivers, container runtimes, schedulers, monitoring systems, and training frameworks all participate in the achieved utilization number. A driver/toolkit incompatibility can silently degrade performance, while a scheduler or monitoring gap can strand accelerators behind failed jobs. Staffing depends on automation, service commitments, and operating scope because performance optimization is continuous: as models, frameworks, and hardware configurations change, communication patterns, memory pressure, and parallelism strategies change with them. Section B.6.1 works an underutilized fleet through the C^3 diagnosis, ruling out communication before isolating the compute axis, which shows the performance engineer how to locate where the lost utilization goes rather than tuning blindly.

↪ The C^3 telemetry diagnostic decision tree is developed in section B.10.

Cloud providers charge a margin above raw hardware and electricity partly because they absorb this operating burden. The point is not that cloud is simpler in every technical sense; it exposes a different interface. The platform either pays for internal expertise in GPU kernel behavior, InfiniBand operations, distributed-systems debugging, liquid cooling, and power infrastructure, or it pays a provider to hide some of that machinery behind a service contract. In both cases, the economics are still governed by useful work per dollar.

12.9.2.2 From TCO to total value of ownership

A more complete framework than TCO is total value of ownership, which includes the value generated by infrastructure rather than only its cost. Two clusters with identical TCO may create different outcomes if one reaches a result two weeks earlier, sustains higher scaling efficiency, starts jobs faster, checkpoints more cheaply, or serves the resulting model at lower cost per token. The value terms are hard to price exactly, but they are not optional in systems reasoning. Time-to-result affects product deadlines. Scaling efficiency affects final model quality under a fixed training window.

Experimentation velocity affects the number of hypotheses the team can test. Inference efficiency can dominate training cost over a multi-year deployment.

This value-oriented perspective changes the infrastructure decision from cost minimization to constraint management. A cheaper cluster that slows the research cycle or produces a model with higher serving cost may be more expensive over the model lifecycle. Conversely, a more expensive training platform can pay for itself if it enables a smaller or more efficient deployed model.

The inference dimension of total value of ownership (TVO) connects a training investment to the service it supports. Section 10.4 establishes why recurring serving cost overtakes a one-time training bill (principle 17); the fleet accounting asks the converse question, whether additional training spend repays itself through cheaper serving. Consider a service of the 175B model that handles 10 million queries per day, each generating an average of 500 tokens, for 5 billion tokens daily. At an assumed $2 per million tokens, serving costs approximately $3.6M per year, held constant over a three-year deployment. If an additional $2M training investment produces a validated 20 percent reduction in serving cost at matched quality and traffic, the three-year gross savings are $2.2M, enough to recover that investment before any additional maintenance costs. A lifecycle forecast must still add retraining and changing demand when they occur.

12.9.3 Depreciation and lifecycle

A three-year-old accelerator fleet that still passes every hardware diagnostic can become uneconomic for frontier pretraining when the next GPU generation delivers 3–4× the performance per watt. It can also lose tens of millions of dollars in hardware value while sitting idle between two annual training runs. ML accelerators depreciate as newer devices change the cost of useful work: a physically functioning accelerator can become uneconomic for a particular training workload while remaining useful for another. The peak arithmetic efficiencies in table 2.2 provide one screening measure, but depreciation and refresh decisions require sustained throughput at matched quality, memory capacity, energy consumption, and the cost of replacing otherwise usable equipment.

Resale value depends on supply, demand, memory capacity, software support, and the workload a buyer intends to run. A refresh budget should therefore test a range of resale assumptions rather than extrapolating a fixed historical depreciation rate across generations.

Refresh policy connects lifecycle planning to scheduling. The three-year horizon used in these examples is a planning assumption, not a universal accounting life or a required replacement date. Replacing the entire fleet at once creates a large CapEx spike, while staggered refresh spreads expenditure but leaves the scheduler with mixed-generation hardware. Placement must then account for different compute capabilities, memory capacities, and communication bandwidths. Keeping generations in separate job pools is one option; jobs that span generations need explicit support for their performance differences.

The same lifecycle logic determines where older accelerators remain useful. A GPU that is inefficient for the largest training runs may still serve smaller models, fine-tuning jobs, or memory-bound inference economically if its capacity, bandwidth, and energy cost meet those workloads' requirements. A newer device's peak arithmetic advantage alone does not settle that comparison. Resale, repurposing, and capacity-sharing arrangements preserve value only when the platform can route useful work to the appropriate generation without hiding performance cliffs from users.

Depreciation is most punitive for bursty workloads. If a team trains one large model per year, the GPUs lose value during the idle months even though they perform no useful computation. Cloud capacity amortizes the same depreciation across many customers, while owned capacity must amortize it across the owner's own workload mix. That is why the utilization invariant appears again: the hardware must either stay busy or be cheap enough to strand.

Isolating hardware cost shows how much depreciation alone charges each run of the 175B model. Under the $350,000-per-8-GPU-node assumption, the 1,024-GPU (128-node) fleet costs $44.8M before network and facility costs. Assume a resale value of $7M–10M after three years in a market where newer accelerators offer 3–4× the performance per watt. That performance assumption does not determine resale value by itself. The assumed residual value yields net hardware depreciation of $34.8M–37.8M. Allocated to only two large training runs, it contributes $17.4M–18.9M per run before electricity, staffing, or facility costs. At 80 percent useful utilization over three years, it instead contributes approximately $1.62–1.76 per used GPU-hour. This is the hardware component of owned

cost, not a complete rate directly comparable with cloud rental. Utilization spreads that commitment across output; the full cost accounting determines economic viability.

12.9.3.1 Refresh screening

Each accelerator generation has raised peak arithmetic per watt, more slowly than raw throughput because device power rose as well. A refresh pays only when the workload realizes enough of the higher ceiling to lower energy and cost per completed run at the same quality target. Memory-bound serving and communication-bound training can therefore justify different refresh choices for the same fleet.

↩ Accelerator generations, their peak rates, the power each draws, and their peak arithmetic per watt are established in section 2.2.3.

A refresh comparison can hold either the arithmetic ceiling or the accelerator power budget fixed, and the two framings answer different questions. Consider replacing a V100 fleet of the reference cluster's size, 1,024 devices with a combined accelerator thermal design power (TDP) of 307.2 kW. On the dense FP16 basis of table 2.2, 130 H100s match that fleet's arithmetic ceiling at about 91 kW of combined TDP. Held at the same 307.2 kW, the budget instead accommodates 438 H100s with 3.4× the original ceiling. The first comparison screens for a lower power envelope at matched peak arithmetic; the second screens for more arithmetic capacity within a fixed power envelope. Neither establishes achieved throughput or measured power.

The fleet matched by arithmetic ceiling is not automatically usable. The original fleet has 7.9× as many devices, so, with the same port count per device, the replacement offers proportionally fewer network ports, and its aggregate HBM falls from 35.2 TB to 11.2 TB. A replacement plan must pass memory-fit, topology, and latency checks before the arithmetic equivalence counts as capacity, and both framings compare peak ceilings rather than energy per completed run at matched quality. The hybrid split described in section 12.9.2 also hedges refresh timing, because cloud bursts provide access to newer hardware before a large purchase commits the fleet to it.

The screening numbers bound the refresh decision for the 175B model without settling it. The H100's 700 W TDP is more than twice the V100's 300 W, yet run energy can still fall, because energy integrates measured power over elapsed time and the faster device finishes sooner. The dollar saving and payback period therefore come from a matched training configuration: measured or validated runtime, accelerator and host power, fabric and storage consumption, and the facility overhead applied to that IT total.

12.9.4 Cloud capacity as an operations constraint

Renting 1,024 cloud GPUs avoids upfront capital expenditure, but a synchronous training job whose tensor-parallel ranks straddle a cloud placement-group boundary stalls on inter-instance network latency at every layer. A spot preemption without checkpoint discipline can erase a 70 percent price discount in a single lost afternoon. For organizations that choose the cloud path, the provider absorbs hardware ownership while exposing capacity, topology, and interruption risk as operating constraints. Instance types and pricing change frequently, so the durable distinction is not vendor branding; it is how the cloud allocation affects communication, scheduling, and fault tolerance.

Three properties of a cloud allocation matter most. The interconnect fabric decides whether distributed jobs keep the scaling assumptions used in design, through fixed-latency and per-byte terms that section D.1 separates. Provisioning granularity decides how much topology the platform controls, from individual instances it must assemble itself to preconnected pods that limit composition. Custom accelerators repeat the generality tax (principle 1): supported operations get cheaper, but end-to-end savings depend on workload mapping, software support, and migration effort. The on-demand, reserved, and spot pricing compared in section 12.9.2 trade optionality, commitment, and interruption risk alongside those physical properties, so the purchasing choice is also a fault-tolerance decision.

Section 8.8.2 develops spot economics from the scheduler's side, where equation 8.3 inflates the spot price by the ratio of total to productive time. The operations view applies that relation to a capacity budget. At a 70 percent discount, a scenario choice within the 60–90 percent range that section cites, spot H100 instances cost approximately \$1.20/GPU-hour instead of \$4. For the 175B model, the 1,024-GPU, 2–4 week workload consumes approximately 344,000–688,000 allocated GPU-hours. The on-demand cost is therefore about \$1.4M–2.8M, while spot pricing would reduce the bill to roughly \$0.4M–0.8M before interruptions. Interruptions, restarts, and lost work can erase that

discount or push the job past its deadline, so completion cost and elapsed time both belong in the comparison. Checkpoint frequency, instance diversification, and on-demand fallback improve that comparison within the limits that section 8.8.3 sets out, and serving can use interruptible capacity only when remaining replicas, routing, and replacement readiness preserve its service contract.

Renting capacity for the 175B model also requires an allocation whose placement contract matches the parallelism plan. Cloud instances share the provider's fabric with other tenants, and placement groups or dedicated fabrics reserve high-bandwidth connectivity only within a group, so a synchronous job whose ranks straddle a group boundary pays that boundary's latency and bandwidth on every step. The working rule is to keep tensor- and pipeline-parallel groups, which exchange data within every layer or micro-batch, inside one placement group, and let only the less frequent data-parallel gradient traffic cross the boundary. Representative collectives must be benchmarked on the actual allocation before nominal GPU availability counts as usable training capacity.

Checkpoint 12.5: TCO decision framework

An organization needs to train 10 models/year, each requiring 1,000 GPU-hours on H100s, evaluating an on-premises cluster with 8 GPUs against cloud instances at $4/GPU-hour.

- ☐ **Cloud cost estimate**: calculate the annual cloud cost from the training volume and hourly rate.
- ☐ **On-premises TCO**: calculate the annualized on-premises cost, assuming $350,000 per 8-GPU node, $0.07/kWh electricity at PUE 1.1, and 700 W per GPU.
- ☐ **Break-even volume**: solve for the number of annual training runs at which on-premises becomes cheaper.

12.9.5 Capacity planning as a reliability control

Sizing a training cluster from parameter count alone can starve the job of memory and fabric bandwidth. Ordering 1,024 H100s that arrive six months before the substation and liquid-cooled racks can power them leaves tens of millions of dollars of silicon sitting in crates while the launch window closes. Infrastructure planning begins when a workload requirement becomes a facility requirement. Model architecture, precision, optimizer state, and batch configuration establish memory needs; the training algorithm and token budget establish work; the deadline establishes required sustained throughput. Parallelization then determines the communication that the fabric must support. Accelerator TDP contributes to the node power envelope, which also includes hosts, memory, and networking, and that envelope constrains rack density and cooling. Planning follows this causal chain from useful work to physical capacity and total cost of ownership.

The first step is workload characterization. Memory sizing follows activations, optimizer state, temporary buffers, and the sharding or offload policy rather than parameter count alone. The token budget fixes total work, and the deadline fixes sustained throughput. That budget is a compute-optimal choice under the scaling laws of section 1.3 (principle 9), and the choice holds only within the breakdown conditions of section 1.3.5. Bottom-up sizing then works from the accelerator outward, using roofline analysis to tell whether each training or serving phase binds on computation or memory bandwidth and the parallelization plan to set the AllReduce, AllToAll, or point-to-point traffic the fabric must carry.

Each layer of physical sizing then applies a method taught in an earlier chapter. Per-device training state sets node memory and the sharding or offload policy (section 2.5.3, section 6.3.3). The primary collective sets the fabric topology and its congestion-control requirements (section 3.5.2, section 3.6.3). Node power sets rack density and cooling (section 2.6.1). The four-node DGX H100 rack assumed for the 175B plan is roughly a 33.5 kW design point once host, network, conversion, and cooling overhead are included. TCO and schedule remain design variables because the cluster must arrive in time to matter, and GPU supply or electrical work often dominates lead time.

↩ The Roofline Model and its ridge point are established in section 2.3.2.

For the 175B training plan, the pod-sizing method of section 2.7 turns the run's 3.15×10^{23} FLOPs budget into a cluster. On a BF16/FP16 planning basis rather than the FP8 headline peak, each H100 sustains an assumed 445.1 TFLOP/s at 45 percent model FLOPs utilization (MFU), so 1,024 GPUs complete the work in about 8 days. That token budget is the GPT-3 operating point, which

is undertrained relative to compute-optimal allocation (section 1.3.2). A plan that followed the compute-optimal ratio of tokens to parameters at the same model size would multiply the token count, and with it the cluster-days, so the scaling-law choice moves the capacity plan before any hardware question arises. A 2–4-week calendar allowance and a 5–10 percent spares pool then cover delays and recovery outside the MFU measurement window; stalls already reflected in MFU must not be counted again.

The plan's 32 racks draw approximately 1.1 MW under the rack accounting boundary of section 2.6.1, and its three-year cost and cloud alternative are those of the running cluster in section 12.9.1, where ownership pays only if utilization stays high beyond the initial training run. Procurement and facility schedules must be confirmed against available capacity. A three-month initial deployment is feasible only if its hardware and facility resources are already available within that window; it cannot be inferred from a longer procurement lead time.

Checkpoint 12.6: Infrastructure planning exercise

A team needs to train a 70B-parameter model on 1T tokens within 4 weeks, with H100 GPUs at 1979 TFLOP/s FP8 peak (45 percent MFU), the headline basis that the worked plan above set aside, and a 2 MW facility power ceiling.

- ☐ **GPU sizing**: determine the minimum number of GPUs needed to complete training within 4 weeks for a compute budget of $6 \times 70 \times 10^9 \times 10^{12} = 4.2 \times 10^{23}$ FLOPs.
- ☐ **Power headroom**: evaluate sufficiency of the 2 MW power budget, assuming PUE 1.1, 700 W per GPU, and 50 percent overhead for non-GPU components.
- ☐ **Cost comparison**: estimate the training-run cost at $4/GPU-hour in the cloud vs. $350,000 per 8-GPU node on premises.
- ☐ **Precision basis**: repeat the GPU sizing on the BF16/FP16 planning peak of 989 TFLOP/s used in the worked plan, explain why the two GPU counts differ, and state which basis a procurement plan should use unless the training recipe has been validated in FP8.

12.9.5.1 Capacity lead time as operational risk

The economics of ML infrastructure are not purely about hardware specifications and electricity rates. Capacity lead time is itself an operational risk: a platform that cannot acquire, reserve, or free capacity quickly enough may miss model launches, delay rollback-safe migrations, or run without the headroom needed for incident response.

The GPU supply chain adds a correlated risk, because processors, memory, packaging, and network equipment that must arrive together can depend on the same upstream capacity, so orders placed with several vendors do not by themselves establish independent supply.

Capacity commitments can take much longer to change than model traffic, which can shift within days of a launch or product change. For planning, consider a large deployment (1,000+ GPUs) with a 6–12-month interval from purchase decision to usable capacity; the actual interval requires supplier and facility confirmation. That delay changes operations. The platform must forecast model growth, reserve refresh capacity, and maintain fallbacks before demand arrives rather than after dashboards turn red.

Facility readiness can lag silicon delivery. In a planning scenario, substation, building-shell, and facility schedules of 12–30 months run concurrently with 6–12-month GPU deliveries, and the longest prerequisite sets the critical path. Site power, data-residency rules, and connection queues can also rule out an otherwise economical location before any accelerator is purchased. The physical design of power delivery and cooling belongs to section 2.6.1; the operational point is that its schedule becomes the fleet's schedule.

Mismatched schedules are expensive in both directions. Hardware delivered before the facility is ready ties up capital without producing useful work, as does a completed facility waiting for accelerators. Phased deployment can reduce the mismatch by commissioning usable capacity in waves while construction continues. An early phase lets the software team validate the distributed-training stack and tune collective kernels on the deployed topology while operators test power

distribution and cooling; success at that scale does not establish readiness at full load. The economic value of making useful capacity available earlier can justify premiums for modular halls or temporary colocation when the expected benefit exceeds their additional cost. Projected monthly model value is an input to that comparison, not a guaranteed loss for every month of delay: the estimate must account for launch readiness, demand, operating costs, and whether delayed value can be recovered later.

Hardware access becomes part of the deployment schedule for the 175B model as well. An illustrative 1,024-H100 allocation represents approximately $45 million in node hardware under the $350,000-per-8-GPU-node assumption established in section 12.9.2, but the operational risk extends beyond the purchase price alone. If the allocation arrives a quarter late, the platform may lack capacity for fine-tuning, shadow traffic, canary expansion, or rollback-safe serving during the model launch window. Capacity planning is therefore a reliability control rather than a procurement function alone.

12.9.6 ML systems TCO framework

A budget that prices only visible GPU compute and object storage misses the engineering labor, pipeline maintenance, continuous retraining, monitoring, and incident response that routinely make up the other half of total cost. Fleet cost accounting prices physical capacity, but deciding which workloads deserve to consume that capacity requires a TCO framework covering the full ML lifecycle alongside the value each workload creates. The lifecycle definition of TCO that opened the fleet analysis (definition 12.1) now applies workload by workload, widening the hardware lens to training and inference costs, with data and iteration completing a four-component model. Their relative weights depend on workload, traffic, retraining cadence, and operating requirements rather than a fixed distinction between ML and traditional software.

12.9.6.1 The TCO equation

Equation 12.12 expresses the total cost of ownership as the sum of four distinct cost components, each with different scaling characteristics and optimization levers:

$$\text{TCO}_{\text{ML}} = C_{\text{train}} + C_{\text{infer}} + C_{\text{total}} + C_{\text{iter}} \tag{12.12}$$

Here, C_{train} is the initial training and evaluation cost, C_{infer} is recurring serving cost, C_{total} covers data acquisition, storage, cleaning, and governance, and C_{iter} captures subsequent experimentation and maintenance. The categories require a common accounting horizon and an explicit allocation rule. If retraining compute is charged to C_{train}, it must be excluded from C_{iter}; shared hardware depreciation and energy likewise enter these workload costs once rather than being added again as a separate fleet bill.

In the illustrative budget of figure 12.14, engineering labor (20 percent), pipeline maintenance (10 percent), continuous retraining, monitoring, incident response, and governance together equal the visible infrastructure spend. The example shows what an incomplete budget can miss; its proportions are not universal cost shares. Each organization must measure these obligations within the same accounting horizon as its infrastructure costs.

The decomposition shows why optimization must follow the measured cost structure. An effort centered on experimentation may spend most on C_{iter}, a training campaign on C_{train}, and a high-traffic deployment on C_{infer}. These are workload regimes, not mandatory stages of organizational maturity. The cost breakdown identifies which component merits attention under the current workload mix.

12.9.6.2 Training cost model

Training cost begins with the price of the GPU-hour. A 30-day run on a 256-node H100 cluster (2,048 GPUs) at the internal fleet rate of $2/GPU-hour is a direct investment of roughly $2.95M. At this scale the cost per training run (C_{train}) becomes a design lever that sizes the fleet. Faster time-to-market pushes the planner toward more GPUs, but larger jobs also suffer diminishing parallel efficiency and a higher probability of hardware interruption. Total cost is therefore not merely a function of compute time; it must account for data staging, checkpointing, and the cost of recovery from interruptions.

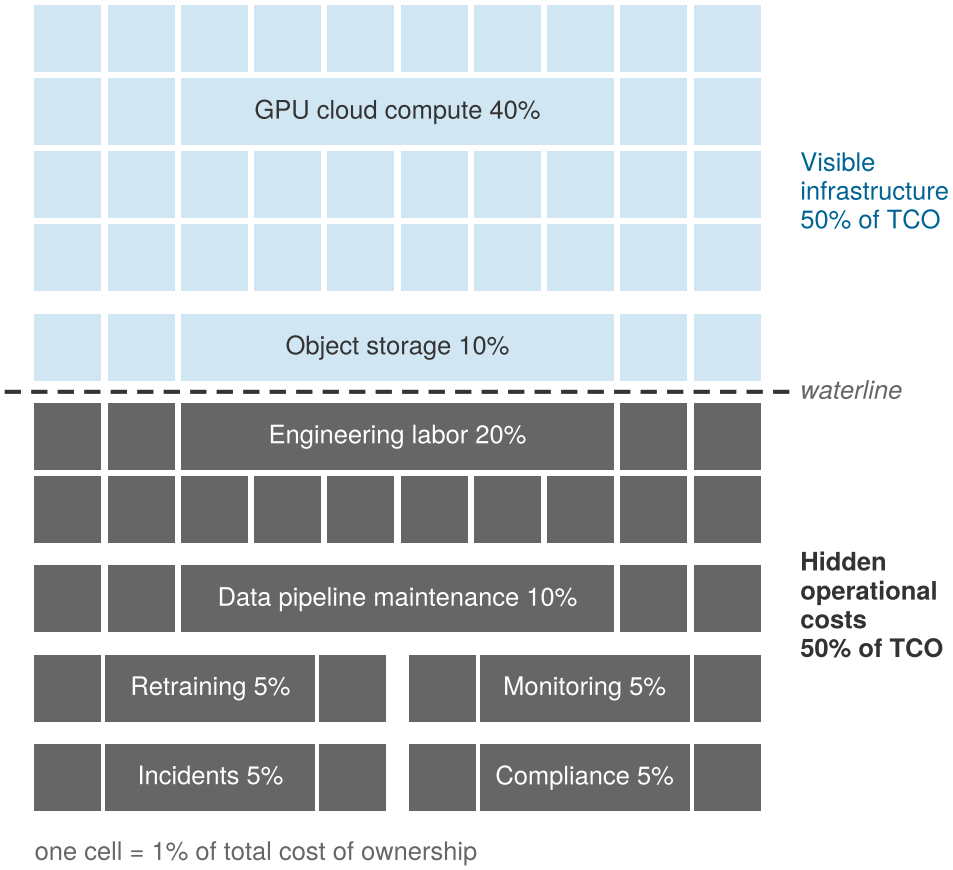

Figure 12.14: The TCO Iceberg: The analysis separates visible infrastructure costs from operational costs that are easier to omit. In this illustrative budget, the hidden operational costs (engineering labor, pipeline maintenance, retraining, monitoring, incident response, and compliance) together constitute half of total cost.

Organizations must decide when to buy additional permanent capacity and when to recover useful throughput through better orchestration or model optimization. The decision also has an energy dimension. Less electricity consumption reduces the energy bill at a fixed tariff, but emissions depend on the energy source and time of use; a lower price alone does not establish lower carbon impact. The cost model below accounts for purchased energy, while an emissions comparison requires the corresponding carbon intensity.

Equation 12.13 formalizes training cost as a function of GPU count, training duration, non-energy accelerator-hour rates, data center energy overhead, and failure overhead:

$$C_{\text{train}} = N_{\text{GPU}} \times T_{\text{hours}} \times \left(R_{\text{GPU/hr}} + P_{\text{device}} \times R_{\$/\text{kWh}} \times \text{PUE}\right) \times (1 + F_{\text{fail}}) \tag{12.13}$$

where:

- N_{GPU} is the number of GPUs allocated
- T_{hours} is the baseline training duration in hours before the additional failure overhead
- $R_{\text{GPU/hr}}$ is the non-energy accelerator-hour cost
- P_{device} is per-GPU power draw in kW
- $R_{\$/\text{kWh}}$ is the electricity price
- PUE is the power usage effectiveness multiplier applied to the energy term

- F_{fail} is additional time from failures and recovery as a fraction of the baseline duration

The PUE factor captures facility energy overhead, the facility metric of section 15.4.1.3. A PUE of 1.1 means 10 percent additional energy relative to IT energy. It applies to energy consumption, not to the whole accelerator-hour rate. This equation uses GPU power as a simplified IT boundary; a full facility estimate must also include host, network, and storage power. The failure overhead F_{fail} adds recovery time beyond the baseline. Checkpointing or restart time already included in measured baseline duration must not be charged again through this factor.

↩ Cluster mean time between failures and the checkpoint interval it implies are established in section 7.1.

12.9.6.3 Training cost calculation

Consider training a large language model:

- Configuration: 256 H100 GPUs for 14 days
- Base non-energy accelerator-hour rate (scenario choice): \$3.50/GPU-hour
- Energy adder: 700 W at \$0.07/kWh and PUE 1.1 adds \$0.05/GPU-hour
- Effective rate: \$3.55/GPU-hour
- Failure overhead: 10 percent (scenario assumption relative to baseline duration)

$C_{\text{train}} = 256 \times (14 \times 24) \times (\$3.50 + \$0.05) \times 1.1 \approx \$335,892$

If this model requires quarterly retraining, annual training cost reaches approximately \$1.3M. However, training cost often represents a small fraction of total TCO for production systems serving millions of users.

12.9.6.4 Inference cost model

Inference costs can dominate when cumulative serving demand outweighs training and iteration costs. Equation 12.14 estimates continuously provisioned capacity from average daily demand and sustained per-GPU throughput measured under representative batching, the target latency SLO, and operating headroom. Daily volume alone does not size traffic peaks; the estimate requires a separate check against the arrival pattern, workload mix, and scaling response:

Serving can outgrow the training bill.

$$N_{\text{GPU}} = \left\lceil \frac{Q_{\text{day}}}{86{,}400 \times R_{\text{GPU,target}}} \right\rceil, \qquad C_{\text{infer}} = N_{\text{GPU}} \times 24 \times 365 \times R_{\text{GPU/hr}} \tag{12.14}$$

where:

- Q_{day} is the daily query volume
- $R_{\text{GPU,target}}$ is sustained queries per second per GPU at the target latency SLO, including batching and operating headroom
- N_{GPU} is the integer number of continuously provisioned GPUs
- $R_{\text{GPU/hr}}$ is the hourly GPU cost

Inference cost therefore grows with query volume in discrete capacity steps. Utilization and batching reduce cost only when they improve measured throughput at the target latency SLO; they should not be applied again as independent divisors.

For capacity planning, the same continuous-provisioning bill can be expressed in terms of allocated GPU-seconds per query:

$$C_{\text{infer}} = Q_{\text{annual}} \times S_{\text{GPU/query}} \times \frac{R_{\text{GPU/hr}}}{3600}$$

where $S_{\text{GPU/query}} = (N_{\text{GPU}} \times 86{,}400)/Q_{\text{day}}$ divides the total GPU-seconds provisioned across the day by the completed queries over the same accounting period. Because the numerator counts every provisioned GPU-second, including off-peak traffic troughs and latency-SLO headroom, substituting $Q_{\text{annual}} = 365 \times Q_{\text{day}}$ into this expression reproduces the continuous-provisioning bill in equation 12.14. Plugging in profiler-measured active kernel time per query instead measures raw device work and excludes idle headroom and diurnal overprovisioning, which substantially understates the serving bill unless a multi-tenant scheduler bills only for active execution slices.

12.9.6.5 Data cost model

Data costs can grow faster than the user base when per-user history, transfer volume, and feature processing all rise together, even though each term below is linear in its own volume. Equation 12.15 decomposes data cost into three components that scale differently: storage cost C_{store}, egress cost C_{egress} for moving data out over the network, and processing cost C_{process}:

$$C_{\text{total}} = C_{\text{store}} + C_{\text{egress}} + C_{\text{process}} \tag{12.15}$$

Storage costs grow with data retention requirements:

$$C_{\text{store}} = D_{\text{vol,store}} \times c_{\text{storage}} \times T_{\text{months}}$$

where $D_{\text{vol,store}}$ is a fixed data volume in GB, c_{storage} is the monthly storage rate, and T_{months} is its retention period in months. This product prices one volume retained at a constant rate; a growing dataset requires accounting for each period's stored volume rather than multiplying the final volume by the entire history.

Egress costs scale with data transfer volume:

$$C_{\text{egress}} = D_{\text{vol,transfer}} \times R_{\text{egress/GB}}$$

At the scenario egress rate of \$0.09/GB, repeated transfers can become a significant cost driver for multi-region deployments and training data distribution. The applicable rate depends on the source, destination, and service contract; the estimate must distinguish billable transfers from included traffic.

Processing costs scale with compute requirements for extract, transform, load, feature engineering, and data validation:

$$C_{\text{process}} = D_{\text{vol,processed}} \times R_{\text{process/GB}}$$

Data processing costs often surprise organizations: a feature engineering pipeline processing 10 TB daily at \$0.02/GB costs \$73,000 annually.

12.9.6.6 Iteration cost model

Development costs capture the engineering investment in experimentation and model improvement. Equation 12.16 formalizes this as the experiment count multiplied by the combined per-experiment engineering and compute costs:

$$C_{\text{iter}} = N_{\text{exp}} \times (C_{\text{engineer,per-exp}} + C_{\text{compute,per-exp}}) \tag{12.16}$$

where:

- N_{exp} is the number of experiments conducted
- $C_{\text{engineer,per-exp}}$ is the engineering cost per experiment
- $C_{\text{compute,per-exp}}$ is the compute cost per experiment

Experiments that do not improve production metrics still incur costs. If 90 percent produce no improvement and every experiment has the same cost, expenditure per improving experiment is $10\times$ the per-experiment cost. A negative result can nevertheless rule out an unproductive direction or expose a risk. Experiment infrastructure should reduce avoidable engineering and compute work while preserving that learning value, rather than optimize the fraction of positive results alone.

12.9.6.7 Startup versus production company TCO

Two constructed organizational scenarios price the same four cost components (table 12.31). The startup operates 1 production model serving 100,000 daily users with monthly retraining, a 2-engineer team, and cloud-native infrastructure. The production company operates 50 models serving 10 million daily users with weekly retraining for high-velocity models, a dedicated 15-engineer ML platform team, and hybrid cloud/on-premise infrastructure. Because traffic, model count, cadence, and staffing all change, the comparison illustrates different cost structures rather than isolating the effect of scale alone.

Table 12.31: TCO Comparison: Startup vs. Production Company: In these constructed budgets, iteration costs dominate the startup scenario, while inference is the largest component in the production scenario. The comparison changes traffic, model count, retraining cadence, and staffing together; it does not isolate the effect of scale. The 100× user increase yields only a 20.6× TCO increase because the startup baseline is mostly salaries, even though inference grows faster than users. Data costs grow 160×, faster than users, in a scenario that also stores features for 50 models and retrains weekly.

Cost component	Startup	Production company	Scaling factor
Training	$5,000/month	$150,000/month	30× (more models, larger)
Inference	$2,000/month	$400,000/month	200× (100× users, 50 models)
Data	$500/month	$80,000/month	160× (50 models, weekly retraining)
Iteration	$40,000/month	$350,000/month	8.75× (team size, experiments)
Total TCO	**$47,500/month**	**$980,000/month**	**20.6×**
Dominant Cost	Iteration (84%)	Inference (41%)	

Two lessons survive the constructed budgets. The dominant term flips, from iteration (84 percent) for the startup to inference (41 percent) for the production company, so the optimization target flips with it. Data cost also grows faster than the user count, because the production scenario stores and processes features for 50 models and retrains weekly, so its storage, egress, and processing volumes (equation 12.15) grow with model count and cadence as well as with users. The budgets are assumed, so they show why measured cost structure matters without isolating model count, feature complexity, or automation as the cause of any ratio.

12.9.6.8 TCO sensitivity analysis

Sensitivity analysis starts with the dollar cost breakdown and changes one workload assumption at a time. Each change touches only some cost terms, and the platform must update specific measurements before it can estimate a new total. More daily users, for example, requires new figures for requests per user, peak concurrency, retained history, and SLO headroom (table 12.32). The same increase in user count can have different consequences for a cached recommendation service and a long-context language-model service.

Table 12.32: TCO Sensitivity Drivers: Each workload change maps to the cost terms it affects and the measurements needed to update them. Shared infrastructure, capacity steps, and changing engineering effort can make the response nonlinear; the complete cost breakdown determines the total.

Workload change	Cost terms affected	Measurements needed
More daily users	Inference, data, and capacity operations	Requests per user, peak concurrency, history retained per user, and headroom needed to meet the SLO
More models	Training, inference, data, and iteration	Refresh cadence, traffic per model, shared feature reuse, and engineering effort per release
More queries per user	Inference and request-related data processing	Request mix, batching, cache hit rate, bytes logged, and provisioned capacity at the target SLO
Larger models	Training, inference, and iteration	Live-state memory, sustained throughput, parallelism, and time spent validating each candidate
More frequent retraining	Training, data preparation, and iteration	Cost per run, incremental data processing, evaluation effort, and any measured serving benefit

User scaling changes requests and stored history; model size scaling changes arithmetic and memory demand. The analysis recomputes those components in dollars over the same period, retaining unchanged components and including iteration and platform overhead, then sums the revised components and compares the total with the baseline. The resulting change identifies where an optimization can create enough recurring savings to repay its implementation cost.

12.9.6.9 Decision framework for speed versus efficiency

TCO analysis enables principled decisions about when to optimize for development speed vs. operational efficiency. Equation 12.17 calculates the breakeven point for optimization investments:

$$T_{\text{breakeven}} = \frac{C_{\text{optimization}}}{C_{\text{save,monthly}}} \tag{12.17}$$

where $C_{\text{optimization}}$ is the one-time implementation and validation cost and $C_{\text{save,monthly}}$ is positive monthly savings net of added operating and maintenance costs. This simple payback assumes savings remain constant over the interval; recovered capacity counts as cash savings only when it reduces expenditure or avoids a planned purchase.

The breakeven equation helps compare specific investments once their implementation costs and recurring savings are estimated. The dominant cost component suggests where to look first, but does not rule out a valuable smaller intervention elsewhere. Table 12.33 offers investigation starting points and illustrative policy thresholds; actual investment criteria also depend on service risk, expected workload lifetime, and available engineering capacity.

Table 12.33: Cost-Structure Decision Rules: Measured cost shares guide investigation and name the risk each posture must evaluate. Candidate savings, useful lifetime, and operational risk determine whether an investment satisfies the organization's policy.

Dominant cost structure	Operating context	Optimization posture and risk to evaluate	Payback threshold
Iteration costs dominate ($C_{\text{iter}} > 50\%$ of TCO)	Experiment-heavy workload	Investigate experiment infrastructure and developer tooling that shorten iteration time. Risk: engineering effort that does not repay through faster iteration or realized savings.	Evaluate candidate value; iteration cost share alone does not set a payback bar.
Training costs dominate ($C_{\text{train}} > 40\%$ of TCO)	Repeated training campaigns	Investigate mixed precision, gradient checkpointing, checkpoint-and-resume on spot instances, and architecture changes. Risk: recovery overhead, completion risk, and useful investment lifetime.	An illustrative policy accepts 3–6 months only if the candidate remains useful beyond payback.
Inference costs dominate ($C_{\text{infer}} > 40\%$ of TCO)	High-volume serving	Investigate quantization, batching, caching, and distillation where repeated queries can realize savings. Risk: quality loss, latency changes, and workload lifetime.	An illustrative policy requires 1–3 months; the actual bar depends on candidate lifetime and risk.
Data costs dominate ($C_{\text{total}} > 30\%$ of TCO)	Feature- and history-heavy	Investigate storage tiering, egress reduction, and feature simplification. Risk: retention, freshness, and downstream quality requirements.	Evaluate against retention and freshness obligations; storage share alone does not set a bar.

The payback bars in the table represent an illustrative investment policy, not a consequence of training or serving alone. Repeated queries can make serving savings accumulate quickly, but a model refresh or request-mix change can retire the optimized path before payback. A reusable training pipeline may support a longer horizon, while a one-off training campaign may not. For either workload, the comparison sets net savings over the candidate's expected useful lifetime against implementation cost and risk; a dominant cost category does not guarantee that its next optimization has the highest value.

12.9.6.10 Optimization investment decisions

A production company with an inference bill of \$400K/month evaluates INT8 quantization:

- Implementation cost: \$80,000 (engineering time + validation)
- Expected inference cost reduction: 40 percent
- Monthly savings: \$400,000 $\times$ 0.40 = \$160,000

$T_{\text{breakeven}} = \frac{\$80{,}000}{\$160{,}000/month} = 0.5$ months

Breakeven in about 2 weeks makes this investment attractive if validation confirms acceptable quality and latency and the expected cost reduction can be realized. For comparison, the same company, with a training bill of \$150K/month, evaluates a training optimization:

- Implementation cost: \$80,000
- Expected training cost reduction: 30 percent

- Monthly savings: \$150,000 × 0.30 = \$45,000

$T_{\text{breakeven}} = \frac{\$80{,}000}{\$45{,}000/month} = 1.8$ months

The training optimization has longer payback and smaller absolute savings in this scenario. Those results favor the inference investment if both satisfy their quality and reliability requirements and compete for the same engineering capacity.

12.9.6.11 TCO-driven architecture decisions

TCO analysis should inform architectural choices alongside operational optimization. The build-vs.-buy analysis in section 12.9.2 also applies to individual workloads, accounting for any capacity they share with the rest of the portfolio. Model architecture selection changes the cost breakdown because a model requiring 2× training cost but 0.5× inference cost may have lower lifecycle cost at sufficient serving volume. Retraining frequency adds to C_{train} and C_{total}; it can reduce emergency work charged to C_{iter} only when evaluation and release controls produce timely, useful updates. Monitoring detects the need to investigate drift; retraining alone does not. Feature complexity can likewise increase C_{total} and C_{train}, while its effect on C_{iter} depends on whether better predictions outweigh added integration and maintenance. Knowledge placement in language-model services illustrates the same trade-off.

12.9.6.12 Retrieval-augmented generation vs. fine-tuning

Domain knowledge can live in a retrievable index, in model weights, or in adapter weights, and each placement creates a different operating cost surface. **Retrieval-augmented generation (RAG)** retrieves documents at inference time and places them in the prompt (Lewis et al. 2020). Fine-tuning stores adaptation in model weights or adapter weights: supervised fine-tuning trains on curated input-output examples, while low-rank adaptation (LoRA) stores small adapter-weight deltas (Hu et al. 2021). Both approaches can improve task behavior, but they move cost and failure risk to different terms of equation 12.12. RAG charges knowledge to C_{infer}, through retrieval and the extra prompt tokens every request must prefill, and to C_{total}, through the index it must build and refresh. Fine-tuning charges it to C_{train} and C_{iter}, through curation, training runs, and re-evaluation whenever the knowledge changes. Query volume therefore decides which cost grows faster.

RAG makes knowledge updates possible through the document index without retraining the model. Retrieved passages provide evidence against which generated claims can be checked, but retrieval alone does not ensure that an answer uses or cites that evidence correctly. Retrieved context also consumes serving resources, through prefill work that grows faster than linearly with sequence length and KV-cache memory that grows with it (section 10.7.1). The cost comparison therefore depends on retrieval quality, context selection, and the serving configuration; excessive context can also make relevant evidence harder for the model to use.

Fine-tuning moves adaptation work before deployment. The platform pays for data curation, training, and evaluation; successful adaptation can shorten prompts when it replaces instructions or examples that would otherwise accompany each request. Lower serving cost is a benefit to measure, not an automatic consequence of changing weights. Adapter-based fine-tuning also creates a fleet-management obligation when many customer-specific LoRA adapters must be versioned, loaded, and isolated, while changing facts embedded in weights requires a different update process from replacing documents in an index.

The decision depends on knowledge volatility, attribution requirements, latency budget, and adapter-management complexity. RAG provides an update path for changing facts and a source trail for review, subject to retrieval and generation quality. Fine-tuning can adapt domain language or stable behavior, subject to training cost and validation. Neither method implies a fixed context length or query-volume limit. A combined design can fine-tune how the model uses tools and documents while retrieving facts that must remain fresh, and evaluation must establish whether those benefits justify the combined lifecycle cost. Because each of these architecture choices is a forecast of lifecycle cost, once a choice ships the fleet needs controls that measure whether spending follows the forecast and give someone the authority to act when it does not.

12.9.7 FinOps for ML platforms

A cost dashboard can identify a failed training sweep without giving anyone authority to stop it. FinOps closes that gap by connecting fleet cost accounting and anomaly signals to budgets, owners, and intervention policies. The platform, fleet, and workload analyses above decide whether capacity and platform investment are worthwhile; FinOps decides how teams spend that capacity and respond when cost or value departs from the plan.

Definition 12.2: FinOps for ML

FinOps for ML is the practice of treating compute cost as a first-class engineering constraint (measured in real-time per experiment and model, and optimized jointly with model accuracy and latency) rather than accounting for it retrospectively through annual budget reconciliation.

1. **Significance**: ML compute costs scale steeply with experimentation volume. A team running 1,000 GPU-hours/day at the fleet reference rate of $2/GPU-hour spends $2,000/day ($730K/year) on training alone. With per-experiment cost visibility, an engineer can identify that 30 percent of runs are terminated early due to divergence (a reason to investigate the training configuration and failure causes) and that the team's fault-tolerant jobs could run on spot capacity at 60–90 percent discounts, provided recovery and interruption costs preserve the required quality and completion target.
2. **Distinction**: Unlike traditional IT budgeting (which allocates fixed annual compute budgets across departments), FinOps for ML operates at the per-run and per-model level with real-time feedback, enabling decisions like early stopping a $50K training run at step 1,000 when loss curves signal divergence, rather than discovering the waste after the full run completes.
3. **Common pitfall**: A frequent misconception is that FinOps means minimizing compute spend. The goal is useful value across the full lifecycle within quality, latency, and reliability requirements: a $500K training run producing a model serving 10B inferences at $0.0001/query may be far more cost-efficient than a $50K run producing a model that requires 10× the inference compute to achieve the same accuracy at scale.

12.9.7.1 Cost components

ML platform cost breakdown depends on the portfolio and accounting period. In an assumed training-heavy annual budget, training compute takes half of spending and serving a quarter, with storage, network, and platform overhead sharing the rest (table 12.34). A mature, high-traffic serving portfolio can instead be inference-dominated, as the earlier lifecycle example shows.

Table 12.34: Illustrative Training-Heavy Platform Budget: Assumed shares sum to 100 percent over one annual accounting period. Training dominates this portfolio; a serving-heavy portfolio can reverse the ordering. Optimization levers belong on the measured cost driver; these shares are not industry norms.

Cost category	Assumed share	Primary drivers	Optimization lever
Training compute	50%	GPU hours, experiment volume	Spot instances, early stopping
Serving compute	25%	Traffic volume, latency SLOs	Autoscaling, model optimization
Storage	12%	Dataset size, checkpoint frequency	Tiered storage, retention policies
Network	8%	Multi-region, data transfer	Caching, compression
Platform overhead	5%	Team size, tooling	Automation, self-service

12.9.7.2 Cost visibility and attribution

Cost optimization requires visibility into spending and ownership. Attribution policy also shapes incentives. Direct metering charges for measured consumption, but shared services and idle reserves still require allocation rules; measured activity does not uniquely determine every team's cost. Charging only for consumption can encourage under-provisioning when teams optimize their

bill instead of service health. Allocation-based attribution charges for reserved capacity, making commitments visible even when that capacity is idle. Hybrid attribution combines a base allocation charge with a variable charge for excess usage. Each policy must preserve the headroom and shared services on which the reliability contract depends. Attribution itself starts from metadata. Tag-based allocation assigns each resource's cost to a team, service, environment, and cost center, as listing 12.7 shows.

Listing 12.7: Cost Attribution Schema: Resource tagging dimensions and shared-cost distribution policies for allocating ML infrastructure expenses to teams and services.

```yaml
# Resource tagging schema for cost attribution
cost_allocation:
  dimensions:
    - team: "recommendation"      # Organizational owner
    - service: "ranking-model"    # Specific service
    - environment: "production"   # prod/staging/dev
    - model_type: "inference"     # training/inference
    - cost_center: "CC-4521"      # Finance tracking

  shared_cost_distribution:
    # Platform infrastructure costs distributed by usage
    - resource: "shared-gpu-cluster"
      method: "proportional_gpu_hours"
    - resource: "feature-store"
      method: "proportional_query_volume"
    - resource: "monitoring-infrastructure"
      method: "equal_split"
```

Shared infrastructure needs its own accounting policy because no single team owns it. Platforms typically adopt one of three models. Proportional allocation distributes shared cluster expenses according to measured consumption metrics such as GPU-hours, storage capacity, or API query volume. Equal split divides common baseline expenses evenly across consuming teams, which is well suited for fixed foundational tooling. Marginal cost charges tenant teams strictly for the incremental compute and storage their workloads introduce. The choice of allocation policy ultimately determines whether operational dashboards can drive organizational action rather than merely describing past spend.

For serving workloads, **cost per inference** provides a unit economic metric. Equation 12.18 divides total serving cost by inference count over the same accounting period. Comparisons require matched request mix, quality and latency targets, and cost boundaries; a short classification request and a long generated response are not equivalent units of work:

$$\text{Cost per inference} = \frac{\text{Total serving cost}}{\text{Total inferences served}} \tag{12.18}$$

Cost per inference then becomes the bridge from infrastructure telemetry to product decisions. Teams can compare model versions by asking whether an accuracy gain justifies a $2\times$ inference cost, evaluate optimization work by measuring whether quantization reduced cost per inference by 40 percent, and forecast monthly serving cost under projected traffic. Tracking the metric by model, customer segment, and request type reveals which workload deserves optimization effort.

Cost dashboards reuse the attention-routing hierarchy of the reliability views in table 12.21. At the top of the hierarchy, the executive view tracks total ML infrastructure cost, its month-over-month trend, budget against actual, and efficiency metrics such as cost per prediction and cost per active user. When that top-level spend shifts, the portfolio view breaks cost down by domain and model type so the platform team can see which domain is driving the change. The service view exposes per-service cost, cost per inference, GPU utilization efficiency, and budget comparison for the responsible team. Finally, the investigation view adds deployment markers and attribution detail so responders can tie an anomaly to a specific change rather than dismissing it as accounting noise.

Chargeback becomes actionable only when fine-grained metering, the attribution rules above, budget forecasts, and anomaly detection feed continuous cost monitoring across four complementary

signals. Cost per inference tracks unit spending under a declared request mix, whereas cost per active user relates spending to the population served, though cross-service comparisons still require comparable activity and product scope. To connect those unit metrics to hardware and calendar limits, GPU utilization efficiency measures how allocated accelerators are used (distinct from revenue or engagement per accelerator-hour, which also depends on product demand and model quality). Budget burn rate tracks spending velocity against the allocation so owners can investigate projected overruns before the budget is exhausted.

12.9.7.3 Cost anomaly detection

Cost anomaly detection must account for the fleet's allocation and lifecycle behavior because several distinct operational events can masquerade as a simple spending spike. Adding a serving replica creates an immediate cost step when that replica occupies multiple accelerators, whereas training sweeps and repeated recovery attempts accumulate GPU-hours gradually without producing useful artifacts. Zombie models (section 12.3.3) hold onto serving capacity long after their traffic has disappeared. Because traffic spikes, shifting request mixes, and cloud price changes can also raise the bill, detection must connect any cost change to allocation, workload, and lifecycle evidence rather than inferring a cause from dollar magnitude alone. An undetected anomaly keeps accruing charges until someone reviews the bill, so cost anomaly detection must balance sensitivity against false positive rates as deliberately as reliability alerting does.

The foundation of cost monitoring is the statistical process control of section 12.6.4.1 applied to spend. For a cost time series with historical mean μ and standard deviation σ, the Z-score quantifies how unusual a current observation is:

$$Z = \frac{C_{\text{current}} - \mu}{\sigma}$$

where C_{current} is the observed cost for the current period.

Complementing Z-score analysis, percentage change detection captures sudden shifts regardless of historical variance, which makes it useful for catching step-function increases such as a misconfigured autoscaler that doubles instance count overnight:

$$\Delta\% = \frac{C_{\text{current}} - C_{\text{previous}}}{C_{\text{previous}}} \times 100$$

Effective alerting requires calibrating thresholds to balance detection sensitivity against operational burden. Two illustrative configurations are a Z-score threshold, which alerts when $|Z| > 3$, and a percentage change threshold, which alerts when $|\Delta\%| > 50\%$ day over day. Seasonal demand and planned jobs must be represented in the baseline before ordinary variation can be interpreted as unusual spending. Under the calibrated Gaussian null of the statistical-process-control rule, a daily $|Z| > 3$ check raises about one raw false alert per year and a 2-sigma check about 16, and section 12.6.2 shows how such crossings compound across many monitored series. Choosing between thresholds requires measuring detection power and the cost of missed events as well as alert burden.

For percentage-based alerts, false positive rates depend on the underlying volatility of costs. Services with naturally variable demand may require higher thresholds (75 percent or 100 percent) to avoid excessive alerts, while stable baseline services can use tighter thresholds (25 percent or 30 percent).

A threshold crossing identifies an unusual observation under the monitoring baseline, not a confirmed defect or cause. Table 12.35 groups candidate cost anomaly root causes into five classes, each with a distinct investigation path.

Table 12.35: Cost Anomaly Root Cause Categories: Five primary categories of cost anomalies with their characteristic indicators and investigation approaches. Traffic increases show proportional QPS growth, while efficiency regressions pair rising latency with stable traffic.

Category	Indicators	Investigation path
Traffic increase	QPS proportional to cost	Check upstream services, marketing campaigns, viral events

Category	Indicators	Investigation path
Efficiency regression	Cost up, QPS unchanged, latency up	Review recent deployments, model updates, infrastructure changes
Resource leak	Gradual cost growth, utilization stable	Check for orphaned resources, failed cleanup jobs, zombie processes
Pricing change	Cost up, all metrics stable	Verify cloud provider pricing, reserved instance expiration
Configuration error	Step-function cost increase	Audit autoscaling rules, instance types, replica counts

12.9.7.4 Detecting an inference cost spike

A recommendation service typically costs $100 per day for inference compute. The operations team receives an alert: today's cost has reached $250 by end of day.

Historical data shows mean daily cost $\mu = \$100$ with standard deviation $\sigma = \$15$, so the first check is whether the alert is statistically plausible under normal operations: $Z = \frac{\$250-\$100}{\$15} = \frac{\$150}{\$15} = 10$. A Z-score of 10 is far outside the assumed Gaussian baseline, but the team first verifies billing completeness, duplicate records, and pipeline correctness. Query volume is unchanged; the investigation must also check request mix and arrival peaks before ruling out demand. In this scenario, those remain comparable while p99 latency rises from 50 ms to 125 ms. A larger ranking model deployed at 2:00 AM adds features, GPU utilization stays near 95 percent, and measured throughput per GPU falls 60 percent. At matched demand and operating targets, that throughput loss implies 2.5× the capacity requirement. The team confirms the allocation increase and unchanged rates before attributing the rise from $100/day to $250/day to the release. It then evaluates whether the measured quality benefit justifies the cost or calls for a leaner serving configuration.

12.9.7.5 Cost optimization strategies

Cost optimization is a risk-allocation decision, not a catalog of discounts. Interruptible compute pays off only when checkpointing bounds the work lost to preemption, serving autoscaling only within the latency budget, and instance right-sizing only against measured cost per step or per inference. Each optimization lever trades a specific risk for a specific saving.

Interruptible training capacity and its checkpoint discipline were priced in section 12.9.4. Serving allocates risk differently. Model load time, lumpy accelerator capacity steps, and the batching trade-off between throughput and latency budget all constrain how far autoscaling can cut serving cost, and section 10.13 develops those controllers. FinOps asks a narrower question, whether the warm spare capacity and fallback paths that absorb demand while new replicas start are priced as deliberate headroom rather than flagged as waste.

Instance selection then starts from the binding constraint rather than the newest instance type. Memory-bound training with large embedding tables prioritizes GPU memory capacity and bandwidth, and compute-bound dense training maximizes sustained FLOP/s per dollar. Latency-sensitive serving minimizes cold start and single-request latency, while throughput-oriented serving maximizes requests per dollar through batching. Selection is therefore benchmark-driven, comparing cost per training step or cost per inference across instance types instead of assuming larger hardware is automatically more economical.

12.9.7.6 Budget-aware development

FinOps extends beyond infrastructure optimization because cost feedback changes which experiments teams run and which models they ship. Unconstrained experimentation leads to runaway costs, but budget controls should make trade-offs visible before they block useful work. Per-experiment limits cap individual training runs at a cost threshold, team budgets allocate monthly compute budgets with visibility into consumption, and approval workflows require review for experiments that exceed cost thresholds. Controls should make the budget boundary and exception path explicit. An authorized owner may raise a limit for a valuable experiment; a job cannot silently exceed it.

Model selection should explicitly consider cost alongside accuracy, and a cost-quality trade-off analysis exposes diminishing returns (table 12.36). Moving from small to medium model yields

a 3 percentage-point accuracy gain for 10× training cost increase, while medium to large yields only 1 additional percentage point for another 10× cost, a pattern that should inform deployment decisions.

Table 12.36: Cost-Quality Trade-Off Analysis: Diminishing returns in model scaling. Small-to-medium model transition yields a 3 percentage-point accuracy gain for 10× cost increase; medium-to-large yields only 1 additional percentage point for another 10× cost. This pattern demonstrates why explicit cost-quality analysis should inform model selection rather than defaulting to larger architectures.

Model	Accuracy	Training cost	Serving cost/1K	Value judgment
Small	92%	$500	$0.10	Baseline
Medium	95%	$5,000	$0.50	3 percentage points for 10× cost
Large	96%	$50,000	$2.00	Additional 1 percentage point for 10× more

The marginal accuracy gain must be valued in the application context. An additional percentage point can matter greatly if it reduces costly errors on an important population, but aggregate accuracy alone does not establish that value. Comparing error consequences, quality by slice, and lifecycle cost prevents defaulting to the largest available model.

Efficiency metrics should be reviewed alongside model quality to identify spending that warrants investigation. Low utilization can reflect required reserve capacity, and experiments without a deployed model can still resolve useful questions. Each efficiency metric therefore prompts an investigation rather than delivering an automatic verdict of waste (table 12.37).

Table 12.37: Efficiency Metric Action Map: Efficiency metrics are useful only when they identify a decision. The platform reviews these signals alongside quality metrics to find waste that accuracy alone hides.

Metric	Waste signal	Typical action
Cost per accuracy point	Accuracy gains are bought with disproportionate spend	Revisit model size, feature set, or serving target before scaling up.
Experiments per production model	Many runs produce little deployable value	Improve experiment design, stopping rules, and promotion criteria.
GPU utilization	Capacity is reserved but not doing useful work	Right-size instances, improve batching, or profile inefficient code.
Spot utilization rate	Eligible workloads use expensive on-demand capacity	Move fault-tolerant jobs to spot instances with checkpoint support.

Taken together, the three economic levels and the FinOps controls make platform investment accountable to measured lifecycle outcomes. Platform return decides whether shared tooling pays, fleet utilization decides whether capacity pays, workload TCO decides where optimization pays, and attribution with anomaly detection keeps each team answerable for its share. A technically sound shared service still needs an owner for its budget, compatibility promises, and retirement decisions. Organizational design determines whether those responsibilities support model teams or leave every cross-team change waiting in a central queue.

12.10 Organizational Operation and Learning

A shared feature change may pass platform checks yet still require approval from several model owners. The control plane becomes durable only when those technical interfaces align with organizational authority. Team structure determines who can approve the change, stop a rollout, or fund the repair when a shared dependency fails.

The central question is where an organization places invariants. Deployment safety, observability, compliance, cost control, and feature consistency can either be enforced once by a platform team or rediscovered differently by every model team.

12.10.1 Centralized platform team

When eight product teams each build their own deployment scripts, feature pipelines, and alert dashboards, the same stale-feature or rollback outage repeats across the company. A centralized ML platform team eliminates that duplication by building and maintaining shared infrastructure while model teams own model development and behavior. Three organizational patterns trade consistency against local responsiveness (figure 12.15). A shared request queue (Panel A) centralizes prioritization but can become a bottleneck; uniform standards depend on enforced contracts, not the queue itself. Embedded tooling (Panel B) favors local responsiveness while risking duplication. The hybrid pattern (Panel C) places dedicated ML engineers near product teams while retaining a common platform. Self-service interfaces can preserve central standards without routing every routine action through a human queue.

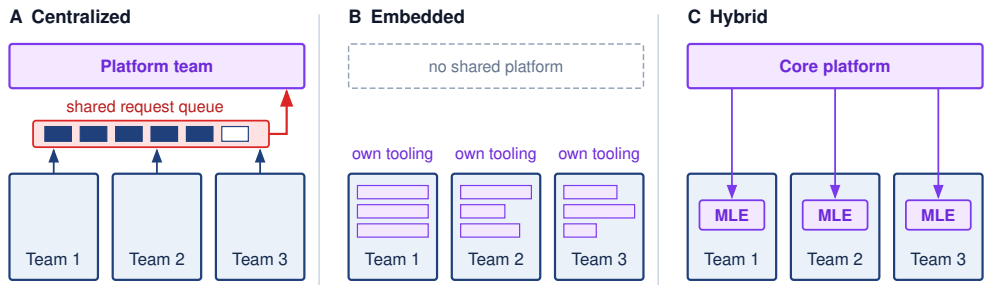

Figure 12.15: Organizational Patterns for ML: (A) A centralized platform team gives every model team the same infrastructure, but all requests pass through one shared queue, which can become the bottleneck. (B) Embedded teams each build their own tooling, which is fast locally but duplicates infrastructure and lets stacks diverge. (C) A hybrid keeps a core platform and places ML engineers (MLE) inside model teams as liaisons, balancing standardization with responsiveness.

Centralization works when the same operational mistakes would otherwise recur across the organization. A shared team can enforce consistent deployment, monitoring, and governance practices; invest once in training, serving, and observability infrastructure; and concentrate deep ML systems expertise that would be hard to replicate across many product teams. It also gives infrastructure engineers a visible career path rather than scattering them as one-person support functions.

The same structure can become a queue. If every platform request routes through one team, model teams wait on prioritization decisions they do not control. Distance from product context can also produce abstractions that are technically elegant but poorly matched to the consuming team's constraints. Centralization is therefore strongest when common invariants dominate local variation, and weakest when model teams need rapid domain-specific changes.

12.10.2 Embedded ML engineers

When a fraud detection team needs a custom streaming feature adapter before tomorrow's attack window, waiting two weeks in a central platform queue is unacceptable. The embedded pattern chooses proximity over uniformity by placing ML infrastructure expertise directly inside model teams, with coordination through communities of practice rather than a single reporting line. The benefit is responsiveness: the infrastructure engineer sees the team's data, model behavior, release pressure, and incident history firsthand, so local changes do not wait for cross-team scheduling.

The cost appears later, when each team has solved deployment, monitoring, and feature management slightly differently. Fragmentation wastes engineering effort, complicates integration, and makes platform quality depend on the skill and bandwidth of whoever happens to be embedded with a given team. Embedded teams therefore work best when local speed matters more than cross-team consistency, and worst when every team is reimplementing the same platform substrate.

12.10.3 Hybrid models

An organization running 80 production models across recommendation, search, fraud, and ads cannot afford four separate serving stacks, nor can one central queue design domain-specific evaluation gates for all four workloads. Hybrid approaches resolve this tension by combining shared infrastructure ownership with domain-specific expertise. In this pattern, the core platform team owns compute, storage, networking, common training and serving systems, monitoring, and shared security, compliance, and cost controls. Domain platform engineers stay close to recommendation, language, vision, fraud, or search teams. Explicit boundaries distinguish changes those teams can make locally from changes that require coordination because they affect shared interfaces or operating guarantees.

A federated version of the same idea distributes implementation work without giving up architectural control. Major contributing teams help maintain platform components, but standards, ownership boundaries, and prioritization are coordinated explicitly. This works only when governance has real authority: otherwise federation degenerates into embedding with more meetings.

12.10.4 Organizational pattern selection

A 15-model startup iterating on disparate prototypes stalls if every change must clear a central governance queue, whereas a 100-model regulated financial institution faces audit failures and repeated outages if every team improvises its own deployment path. The appropriate organizational pattern depends on where standardization creates more value than local autonomy. Higher model counts, stricter regulatory requirements, and immature infrastructure favor centralized platforms, while heterogeneous model portfolios and smaller organizations may benefit from distributed expertise (table 12.38).

Table 12.38: Organizational Pattern Decision Factors: Five criteria for choosing between centralized and distributed ML platform teams. Higher model counts (100+), stricter regulatory requirements, and earlier infrastructure maturity favor centralized platforms; heterogeneous model portfolios and smaller organizations may benefit from distributed expertise with coordination through communities of practice.

Factor	Favors centralized	Favors distributed
Model count	Higher (100+)	Lower (10–20)
Model similarity	Homogeneous	Heterogeneous
Organization size	Larger	Smaller
Regulatory requirements	Stricter	Lighter
Infrastructure maturity	Earlier stage	Later stage

The infrastructure-maturity row can look inverted, because a young organization seems the natural home for local improvisation. The reasoning runs the other way. Early on, no shared substrate exists, and a central team is the cheapest way to build one before every model team writes its own deployment and monitoring path. Once that substrate matures into self-service interfaces, distributed teams can consume it directly, and the central queue matters less.

The selection logic becomes clearer in a concrete scenario. A technology company with 50 ML engineers across 8 teams operates 80 production models across recommendation, fraud, search, and ads. Each team maintains its own deployment and monitoring, creating duplicated work and inconsistent integration practices. That duplication motivates evaluating a shared platform; the model count alone does not establish its return. Domain diversity also motivates retaining expertise close to each workload rather than expecting one central team to make every modeling and operational decision.

One candidate design is hybrid. In this scenario, a central platform team of roughly 12–15 engineers owns core infrastructure, while domain-specific platform leads remain embedded with the largest model teams. That staffing assumption must be tested against the platform's scope, support burden, and service commitments. A community of practice coordinates standards, and a shared contribution model lets domain teams upstream reusable components instead of maintaining permanent forks. The same platform engineering, feature management, and governance trade-offs become more concrete in the hyper-scale systems built by large technology companies.

12.11 Case Studies

When hundreds of engineers train and deploy thousands of models against shared data streams and strict latency budgets, bespoke per-team pipelines collapse under silent feature skew, duplicated glue code, and uncoordinated rollouts. The five platform cases that follow instantiate the operations trade-offs developed earlier in this chapter: standardization vs. team-local velocity, abstraction vs. control, freshness vs. cost, and governance vs. autonomy. Each case shows a different way to keep a multi-model fleet repeatable without turning every model into bespoke infrastructure.

12.11.1 Uber Michelangelo

Uber's Michelangelo platform shows why repeatability becomes a platform requirement at company scale (Hermann and Del Balso 2017). Uber's 2017 account describes data scientists using varied modeling tools while engineering teams built custom production systems, with no common deployment path or reproducible data pipelines. Michelangelo standardized one workflow from data management through training, evaluation, deployment, and prediction monitoring, and shared feature definitions let teams reuse platform-managed pipelines. Its delivery-time model combined historical and near-real-time restaurant preparation features, and such features raise the consistency problem of section 12.4, because one feature definition must produce compatible values for offline training sets and online requests. The standardization trade-off is the durable lesson. Common data and deployment paths remove repeated infrastructure work, while product teams still own model behavior and business-specific validation.

12.11.2 Meta FBLearner Flow

Meta's FBLearner Flow shows how shared ML infrastructure can hide resource management while preserving workflow structure (Engineering 2016). In its 2016 account, Facebook reported use by more than 25 percent of its engineering team and more than one million models trained. Engineers defined operators with typed inputs and outputs, and compilation formed a DAG whose dependencies controlled execution; operators declared CPU, GPU, and memory needs, and the platform allocated resources and moved data between steps. Feed ranking adds freshness pressure, because the value of a signal such as a recent like decays quickly. Shadow mode (or dark launching) and canary releases then verify a candidate's latency, memory, and error behavior on live traffic before its outputs reach users. The durable lesson is the interface between workflow intent and managed execution.

12.11.3 Netflix ML infrastructure

Netflix's ML infrastructure shows how personalization platforms manage fan-out cost under a strict latency budget (Gomez-Uribe and Hunt 2015). The recommendation system combines several ranking and personalization algorithms for page generation, row selection, search, and artwork, and every additional online ranker or feature dependency consumes latency budget. Production systems therefore separate candidate generation, offline training, online ranking, and caching, so the homepage stays personalized without executing an unbounded number of models per request. New members and new titles raise the cold-start problem; exploration policies such as bandits, which route some traffic to options whose value is still uncertain, are one option, and they need explicit exploration budgets, policy versions, and feedback logging. Netflix's interleaving work merges competing ranked lists for the same user and credits each click to its source ranker (Netflix Technology Blog 2017), so faster experiments require additional serving and logging capability.

12.11.4 Google Vertex AI

Google Vertex AI illustrates the managed-platform abstraction problem, which is to reduce the glue-code burden described by Sculley et al. (2015) while preserving control over custom workloads. Offering AutoML, automated search over model and training configurations, alongside custom training makes that boundary visible, because such search, of which neural architecture search is one method, trades compute for engineering effort and needs bounded cost and concurrency. A managed service also carries contracts that its consumers must verify. A container separates a model's software environment but does not isolate accelerator memory bandwidth, shared networks, or backing

services, and a managed feature service must respect when information was available to the predictor, not merely when the event occurred. Managed infrastructure transfers implementation work, not responsibility for leakage, quality, or tail latency under the intended load. These mechanisms are general design options, not claims about Vertex AI's internal architecture.

12.11.5 Spotify ML platform

Spotify's ML platform places shared infrastructure in a personalization setting where immediate response and longer-term listener value need not align, since optimizing clicks alone can overlook diversity and discovery (the concern about filter bubbles). Separating candidate generators, such as collaborative filtering models (Koren et al. 2009), which recommend what similar listeners chose, or editorial selections, from rankers lets teams change retrieval strategies while preserving the ranking contract. Spotify's documented contribution is a paved road for model orchestration built around TensorFlow Extended, Kubeflow Pipelines, and metadata tracking (Engineering 2019). Its operational lesson is lineage, which lets an engineer trace a silent regression back to the data or workflow state that produced it. The same discipline supports mixed freshness. In an illustrative split, a weekly playlist can be computed in batch, while a responsive home screen ranks precomputed candidates with recent interaction features online, so only decisions that need current context pay for the live path.

12.11.6 Comparing platform control planes

Across Uber, Meta, Netflix, Google, and Spotify, teams started from different bottlenecks (inconsistent delivery-time features, duplicated DAG plumbing, homepage fan-out latency, multi-tenant cloud glue code, and unrepeatable personalization experiments) yet converged on the same division of control-plane responsibilities. No single implementation is the template; the durable pattern is that shared state, release evidence, resource policy, and ownership become platform interfaces, while each domain team keeps a responsibility that no platform can absorb (table 12.39).

Table 12.39: Comparative Platform Control Planes: The five cases allocate implementation differently but retain the same systems contract: the platform owns shared invariants and evidence, while domain teams retain workload intent and acceptance criteria.

Case	Primary pressure	Control-plane emphasis	Local responsibility retained
Michelangelo	Repeated feature and deployment work	Shared features and repeatable workflows	Product-specific model behavior and validation
FBLearner Flow	Many teams and heterogeneous pipelines	DAG orchestration and resource abstraction	Workflow intent and domain evaluation
Netflix	Fan-out under strict latency budgets	Composition, precomputation, experimentation	Ranking objectives and user-experience trade-offs
Vertex AI	Multi-tenant managed infrastructure	Isolation, policy, and lifecycle integration	Workload configuration and acceptance criteria
Spotify	Fresh personalization with mixed paths	Lineage, paved roads, and hybrid data movement	Domain policy and exploration strategy

The comparison also exposes the adoption constraint. A platform succeeds when its common path captures recurring invariants without erasing domain judgment. Too little platform ownership recreates bespoke infrastructure; too much converts every exceptional workload into a central-team queue. Versioned interfaces, explicit exceptions, and outcome-based adoption metrics keep that boundary adjustable as the portfolio changes. The same boundary choices, misjudged, produce the misconceptions that recur when teams scale ML operations.

12.12 Fallacies and Pitfalls

Most operational misconceptions at scale share one root. Practices sized for one model survive into a fleet whose models share features, pipelines, infrastructure, and alerts. The following fallacies and pitfalls show where those per-model practices fail once the interactions dominate.

Fallacy: *Operational complexity grows linearly with model count.*

Shared dependencies can make coordination grow faster than model count. A model consuming embeddings from another model through a shared feature pipeline cannot change that interface safely without checking its consumers. The graph's structure, not the count alone, determines the interaction burden; compatible interfaces still permit independent releases. In the monitoring scenario, 100 models with 10 metrics checked every five minutes at a 5 percent false positive rate produce 14,400 raw false crossings daily before suppression, and at an assumed 40 operational engineer-hours per model per month, 40 models require 1,600 engineer-hours monthly ($2.88M annually at $150/hour). Both figures scale linearly with model count, so they mark the floor of the operational load. Above that floor, N_{models} models form roughly $N_{\text{models}}^2/2$ possible pairs (section 12.2), and the recorded dependency edges decide which of those pairs can actually conflict. A staffing plan that adds operators in proportion to model count therefore falls behind exactly when the dependency graph densifies, which is when shared controls must absorb the repeated work and cross-model interactions instead.

Pitfall: *Applying independent alerts to every model and metric.*

Paging on every per-model threshold crossing creates avoidable noise, because equation 12.9 makes at least one false crossing nearly certain across a thousand independent checks, even though that probability says nothing about the fraction of pages that are false. Hierarchical monitoring supplies context for grouping and routing related signals, while persistence rules and calibrated thresholds reduce nonactionable notifications. Urgent model or component failures still require direct alerts when their response contracts demand immediate action.

Fallacy: *Platform investment can wait until the organization reaches 100+ models.*

The benefit-cost curves charted in figure 12.2 reach break-even parity between 20 and 50 models under the example's assumed costs and savings. The governing threshold is recovered engineering capacity relative to the full platform cost, not a universal model count. The debt-audit examples identify recurring work that can justify investment, but a migration premium or payback period requires its own effort and recurring-cost estimates.

Pitfall: *Treating all deployments with uniform procedures regardless of risk profile.*

Applying the same staged rollout to every model update either over-burdens low-risk changes or under-protects high-risk ones. Release cadences differ by workload family (table 12.2), but a model family does not establish a universal release frequency or rollback window. Risk-based policies match observation time, exposure, and containment to the change: a subtle quality regression may require prolonged shadow evaluation, while a harmful release still requires rapid routing to a validated version.

Fallacy: *Per-model metrics, such as loss and accuracy, are sufficient at scale.*

In multi-model platforms, system-level and per-model metrics answer different questions. An upstream latency regression can lengthen the request path while accuracy remains unchanged. Shared-feature drift can reduce several downstream models' accuracy at once; their metrics reveal the harm but do not alone identify its common cause. Section 12.6 joins business, portfolio, model, and infrastructure signals so the platform can trace that dependency.

Pitfall: *Defaulting to batch pipelines for all features to simplify architecture.*

A daily batch policy is appropriate only when its delay fits the feature contract. The freshness definition $T_{\text{freshness}} = t - T_{\text{event}}$ measures event-to-availability delay; it does not determine a universal batch or streaming latency. The scenarios $T_{\text{freshness}} \approx 12$–$24$ hours and $T_{\text{freshness}} \approx 1$–$5$ seconds illustrate different operating points whose quality effects require measurement. For the recommendation example in table 12.10, the 2.1-point gap between hourly and real-time features is assumed to be worth $10 million a year. A batch default forgoes that value silently, so the streaming decision belongs in the freshness break-even of checkpoint 12.2 rather than in a blanket preference for simpler pipelines. Fraud systems likewise need timely evidence, but the exploitable delay depends on their other detection and containment controls.

Fallacy: *Technical debt is inevitable at scale and should be addressed only when it blocks critical work.*

The premise overlooks the accumulating cost of recurring work. Section 12.2.4.1 provides illustrative escalation policies, not universal thresholds for declaring debt. Release friction, incidents, and toil are tracked against the platform's own baseline, and the cost of remediation is compared with its expected benefit. A repeated manual task becomes an investment candidate when reducing it releases enough capacity or risk to justify the work. Waiting for debt to block critical work also means waiting until it is expensive to redirect. Deployment lead time, incident frequency, alert noise, rollback delay, and toil share are leading indicators that show the platform boundary failing while the repair is still small, and dismissing their rise as the price of scale discards that signal.

12.13 Summary

Fleet operations makes shared model and feature state, release evidence, resource commitments, and recovery authority governable across a portfolio. Registries, dependency-aware deployment, and hierarchical monitoring replace repeated per-model coordination with accountable shared controls. Coordination is the organizing concern in C^3, but those controls also consume compute and communication resources. Within one job, stragglers lengthen the Fleet Law's synchronization term $T_{\text{sync}}(N)$, while data stalls, communication delays, and checkpoint recovery cost time through other mechanisms and on other timescales; across the portfolio, ML productivity goodput charges each lost hour to scheduling, runtime, or program goodput, so the portfolio dashboard names the layer, and the owner, that must act. Registries and release gates connect the fleet's resource use to validated, useful outcomes.

Three principles carry forward from those mechanisms. Release evidence is valid only when it binds to one versioned contract, so offline quality, production load, experiment design, and rollback thresholds must describe the same candidate, and any exhausted dimension of the reliability budget vector stops promotion. A regression is priced by exposure as well as by detection, so rollout speed and rollback readiness follow the harm a bad release can do and the time needed to see it. Utilization is the economic invariant beneath every capacity decision, from build-vs.-buy to standby headroom, while the lifecycle cost breakdown ($\text{TCO}_{\text{ML}} = C_{\text{train}} + C_{\text{infer}} + C_{\text{total}} + C_{\text{iter}}$) decides where optimization repays its cost. Edge fleets apply the same contract under weeks-long rollouts, version skew, and privacy-limited telemetry (section 11.10.1.5).

Key Takeaways: Operational load follows coupling

- **One model is not the unit**: At portfolio scale, the operational object becomes the dependency graph of models, features, pipelines, alerts, and owners. Registries, lineage, and ensemble-aware CI/CD expose affected consumers and make compatibility a release condition.
- **Feature meaning is a versioned interface**: A redefined feature can pass schema checks while silently changing every consumer's inputs. Versioned definitions, lineage, and point-in-time joins that filter on availability time keep training and serving on the same meaning.
- **Shared controls must earn their cost**: At fixed savings per model, the platform benefit-cost ratio grows linearly with model count, so the break-even portfolio size of figure 12.2 moves with savings per model and platform cost. Treating the platform as a versioned product keeps adoption, exceptions, migrations, and retirement from recreating that work.
- **Ownership cost chooses the target**: The cost-of-ownership equation separates training, inference, data, and iteration costs, and the dominant term changes over a system's life (principle 17). Measured lifecycle costs, not fleet age, decide between investing in iteration speed and in serving efficiency.
- **Reliability is a vector**: Independent metric alerts scale into noise across a model portfolio. Hierarchical telemetry routes common-cause incidents, while any exhausted availability,

latency, freshness, or quality budget can stop promotion, because $\max_i B_i \leq 1$ governs rather than an average (principle 21).

- **Evaluation is one versioned contract**: Offline quality, production load, statistical power, assurance evidence, rollback thresholds, and reproducibility metadata must describe the same candidate (principle 21). Every acceptance claim needs an owner and an invalidation condition, so the contract turns the coordination axis of C^3 into auditable evidence.

Shared controls pay off when they reduce the coordination required for a safe change without obscuring who owns its consequences. The platform can automate checks, preserve evidence, and enforce limits, but model and domain owners still decide what behavior is acceptable. That boundary makes fleet growth manageable without turning every exception into a central-team bottleneck.

What's Next: A governed fleet is still an attack surface

Even when every release is versioned, gated, and observable, the fleet is not secure by that evidence alone, because reliability metrics show whether the fleet serves its intended callers, not what an adversary can learn or change through the same interfaces. The serving endpoints, shared feature pipelines, telemetry stores, and federated rounds that this chapter made governable are also places where an authenticated caller can copy model behavior or a participant can alter what the fleet learns. Defending them starts from a different inventory than operating them, one that begins with the adversary rather than the service.

↪ Threat models, fleet trust boundaries, model extraction, data poisoning, and differential privacy are developed in chapter 13.

IV

THE RESPONSIBLE FLEET

Part IV

12

Responsible Fleet Principles

Parts I through III built the machine learning fleet, coordinated its training, and optimized its deployment. Part IV develops the mechanisms for enforcing the security, privacy, robustness, environmental, and social constraints that have shaped those design decisions from the outset. Their detailed treatment follows deployment because enforcement depends on the infrastructure, training, and operational controls already introduced. Each of these constraints can be stated as a budget, measured, and exceeded, which makes it an engineering constraint in the same sense as bandwidth, power, or latency.

What these constraints share is that the evidence collected in the earlier parts cannot certify them. A fleet can meet its latency targets and pass every release gate while an authenticated caller copies the model's behavior, an authorized answer reveals a training record, or a poisoned contribution plants a backdoor. Its health signals can stay normal while its predictions degrade under shifted inputs. Its energy per request can fall while its total energy rises, and an accurate, fast model can still concentrate its errors among the people least able to contest them. No single algorithm closes any of these gaps; each requires monitoring, verification, and governance systems that surround the fleet. Ignoring them is not a neutral choice, because a breach, a regulatory action, or a site that cannot supply power can stop a fleet as surely as a hardware failure. The principles below preview the limits that the following chapters make operational. They identify decisions that belong in the initial system requirements even though the mechanisms for evaluating and enforcing them require the full fleet context.

The first boundary is what the fleet can trust. A model's behavior is shaped by far more parties than the team that trains it.

Principle 22: The provenance bound on trust

Invariant: A trained model's behavior is a function of every contributor to its data, code, and weights, including training records and labels, preprocessing and training code, software dependencies, pretrained checkpoints, and the hardware and firmware that executed the run. A contributor that cannot be identified and verified can change that behavior without producing a visible failure, so the provenance record bounds how far the model can be trusted.

Implication: Poisoning and supply-chain attacks need not crash anything. A small set of crafted training examples or a backdoored checkpoint can shift a targeted decision while aggregate validation accuracy stays unchanged, so clean evaluation results cannot certify integrity. Trust must be established upstream, through signed artifacts, immutable lineage for data and code, controlled dependencies, and verification before an artifact enters training or deployment, and it extends only as far as that record reaches.

The second boundary is what the fleet can reveal. Even when every upstream artifact is verified, a deployed model can still leak private training records through the predictions it returns to callers.

Principle 23: The information leakage invariant

Invariant: A privacy claim must name the protected unit and the observations available to an adversary. Let R be one protected individual's training record and O an observable model output. Leakage about that record is present when

$$I(R;O) > 0.$$

Here, $I(\cdot;\cdot)$ denotes mutual information.

Implication: A correctly executing model can disclose its training data through the outputs it is authorized to give, as memorized text, as membership, or as reconstructed inputs, so access control over stored data cannot protect it. Privacy is a budget, not a switch. Post hoc anonymization of input records does not remove sensitive information already memorized in model weights. Systems requiring formal record-level privacy guarantees should use mechanisms such as differential privacy (DP) where appropriate. DP bounds the influence of adding or removing one protected record for a data analysis or training procedure. Interactive query systems must account for composition, and they may stop answering once the system exhausts the allocated privacy budget.

Privacy defines what the fleet may reveal; robustness defines how reliably it behaves when distributions drift, environments change, or inputs actively probe its weaknesses. Hardening a model against those perturbations is not a free post-processing step, because training against worst-case inputs multiplies the arithmetic required for each update.

Principle 24: The robustness compute penalty

Invariant: Adversarial training buys intrinsic robustness with training compute. It trains on adversarial examples, inputs perturbed within a stated budget so that the model's prediction changes. Projected gradient descent (PGD)-style training finds each perturbation with K_{PGD} inner attack steps, so a batch costs roughly $K_{\text{PGD}} + 1$ forward-backward passes instead of one. The resulting clean-accuracy and robustness effects depend on the attack strength and training procedure.

Implication: Robustness is claimed only per stated failure class (distribution shift, adversarial input, poisoning, silent hardware fault), each with its own evidence, because every one of them can keep service-health signals normal while predictions become invalid. Robustness measures consume training or serving resources and must be evaluated against a stated threat model; certified robustness, for example, is paid for in serving samples rather than training passes. External guardrails such as input filtering and output verification can complement adversarial training, but their detection errors and latency costs remain part of the system. They do not establish the same guarantee as robustness of the underlying model.

Efficiency gains can also change demand, so lower energy per request does not by itself establish lower fleet energy consumption. When cheaper inference invites broader deployment, longer contexts, or more test-time search, aggregate power draw can climb even as each operation becomes leaner.

Principle 25: The Jevons paradox of AI (efficiency trap)

Invariant: Efficiency improvements can lower unit cost and stimulate demand. In this comparison, total resource consumption rises only when usage grows by a larger factor than resource efficiency improves. The following schematic depicts that rebound scenario, not an inevitable consequence of efficiency:

$$\text{Efficiency} \uparrow \implies \text{Cost} \downarrow \implies \text{Demand} \uparrow\uparrow$$

Implication: Making models 10× more efficient can increase total usage enough to erase or even exceed the expected energy savings. Rebound is not the only path. Scaling laws price each further loss reduction in multiplicative compute, so demand can outrun efficiency even without a demand response to lower cost. Sustainability strategies must focus on absolute limits (carbon budgets, renewable sourcing) rather than rate efficiency (FLOP/s per watt) alone.

Responsible operation also requires making social objectives explicit. Calibration asks whether predicted probabilities match observed frequencies. Score balance compares average scores across groups separately among positive and negative outcomes; equalized odds instead compares thresholded error rates. These criteria can conflict, so a fairness requirement must specify which guarantee the system should enforce.

Principle 26: The fairness impossibility theorem

Invariant: For imperfect risk scores across groups with different base rates, groupwise calibration and equal mean scores across groups within each outcome class cannot all hold simultaneously.

$$\Pr(Y = 1 \mid G = a) \neq \Pr(Y = 1 \mid G = b) \implies \text{Trade-off Required}$$

Here, Y is the outcome and $G \in \{a, b\}$ is the group attribute; the implication assumes imperfect prediction and refers specifically to the three score criteria above. Equalized odds at one chosen decision threshold is a different condition.

Implication: Under the theorem's conditions, fairness requires an explicit trade-off; there is no criterion-free optimum that selects among incompatible metrics. Engineers must treat fairness metrics like latency budgets: explicit trade-offs chosen by stakeholders, enforced by the system, and monitored for violation. Chapter 16 provides the full treatment of when the impossibility result applies, including the equal-base-rate and perfect-prediction exceptions.

Those trade-offs cannot be evaluated once and frozen, because deployed systems alter the data and incentives they later observe. Once a model's predictions steer user behavior, ranking exposure, or eligibility decisions, the incoming distribution shifts in response to the system itself.

Principle 27: The sociotechnical feedback invariant

Invariant: When a deployed model's decisions affect its environment, the probability distribution of future data $p_{t+1}(x)$ depends on the model's past decisions $f_t(x)$.

$$p_{t+1}(x) = g(p_t(x), f_t(x))$$

Here, g is the environment's response function, which maps the current data distribution and the model's decisions into the next data distribution. Settings with this feedback, where predictions change the outcomes they predict, are called performative.

Implication: Systems whose decisions affect their environments require closed-loop governance. A model that maximizes accuracy on static test data can still degrade the future data distribution it operates on, amplify incentives that bias the next round of data, or destabilize the environment being monitored. Reliability in these settings requires modeling the *feedback loop*, not the *feed-forward inference* alone.

Part IV develops these previews in dependency order. Security and privacy come first, because they decide who may reach the model, what its releases may disclose, and which contributors it may trust; the threat model for data poisoning begins there. Access controls cannot show that the model's answers stay valid, so robustness follows. It asks whether behavior remains acceptable when inputs, environments, hardware, or adversaries depart from validated conditions, and it

develops the evidence and defenses that each failure class requires. Both are paid for in compute, which brings sustainability next. It treats planetary-scale infrastructure as an energy and carbon system whose grid power and facility overhead bound how much of that compute the fleet can run. Responsible AI closes the part by asking whether a fleet that is secure, robust, and within its resource budget is fit to deploy, through an assurance case that an accountable owner can block. It also connects those technical controls to institutional governance. None of these controls is free in C^3 terms. Security and robustness spend compute and serving latency on verification. Sustainability prices every powered second of the Fleet Law's computation, communication, and synchronization terms. Governance is the coordination that decides when the fleet may run at all. Together, these chapters complete the engineering discipline that turns a capable fleet into a trustworthy one.

13

Security & Privacy

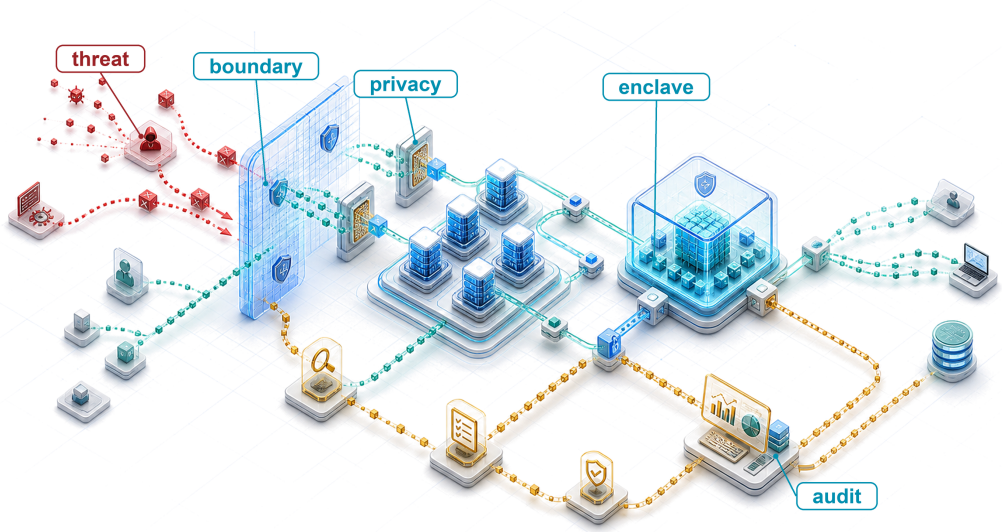

13.1 The Expanded Attack Surface
13.2 Security and Privacy Foundations
13.3 Learning from Security Breaches
13.4 Systematic Threat Analysis and Risk Assessment
13.5 Model-Specific Attack Vectors
13.6 Fleet-Internal Trust Boundaries
13.7 Hardware-Level Security Vulnerabilities
13.8 Comprehensive Defense Architectures
13.9 Differential Privacy
13.10 Selecting and Sequencing Defenses
13.11 Fallacies and Pitfalls
13.12 Summary

Purpose

Why can a correctly functioning machine learning service still disclose sensitive information or enable an attack?

Many useful machine learning systems depend on personal data, institutional knowledge, or behavioral patterns, creating tension between utility and protection. Models can encode sensitive information and later reveal it through their outputs, extending risk beyond transient data processing. An authorized request can expose a private training example, while accepted training data can be manipulated to alter future behavior. In both cases, the software may execute as designed while the service violates its intended boundaries. Protecting a single server therefore leaves questions about what the model learned, who supplied it, and what callers can induce it to do. Scaling introduces more participants, copies, credentials, and handoffs, often across organizations with different responsibilities. A control applied at one boundary may leave another exposed, and a compromised participant can affect work shared across the fleet. The engineering decision is which actors to trust with which information and authority, then how to enforce those choices throughout training, release, and operation. Protection also consumes resources: restricting access can complicate coordination, checking requests adds work, and limiting information release can reduce model utility. Those costs must fit the service's operating budget without making safeguards optional when demand rises. A high-performing model may still be undeployable if sensitive information or control escapes the intended boundary. In C^3 terms, Compute, Communication, and Coordination must preserve confidentiality, integrity, and accountable control as well as throughput.

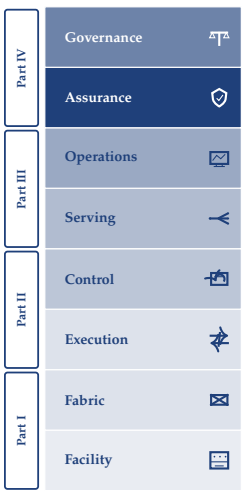

Learning Objectives

- Distinguish security from privacy in ML systems through formal definitions, threat models, and quantitative trade-offs
- Explain how a correctly executing model can disclose memorized training data through authorized outputs
- Analyze ML-specific attack vectors across model theft, data poisoning, adversarial examples, fleet-internal gradient exchange, and hardware vulnerabilities
- Apply lessons from historical breaches and a layered maturity model to prioritize controls for a deployment's threat model
- Design layered defenses across data, model, runtime, and hardware, weighing each control by the trust boundary it protects and the compute, communication, or coordination cost it imposes
- Specify privacy-accounting inputs and assess reported budgets against sampling assumptions, composition, and measured utility
- Select rollback and forensic responses for a suspect model by tracing its provenance through the release record

13.1 The Expanded Attack Surface

A serving fleet can meet its latency target while an authenticated client collects enough responses to copy model behavior. A federated training job can converge while a participating device submits poisoned updates. The release gates and hierarchical monitoring that govern the fleet can report both systems healthy, because a copied model and a planted backdoor need not move any availability, latency, or aggregate-quality signal. Operational health alone therefore cannot establish security. In the fleet stack (figure 1.4), the governance band begins with the Assurance row, which sits directly above Operations and which this chapter shares with chapter 14. That row must preserve the confidentiality, integrity, and availability of the fleet across the same interfaces that make distributed training and serving possible.

Selecting a defense starts by identifying the protected asset, the actor who threatens it, the access that actor has, and the boundary that access must not cross. An API client can observe responses; a training participant can submit updates; a host administrator may inspect runtime state. Those capabilities call for different controls. The breach cases make these boundary failures concrete, the threat model turns them into explicit assumptions, and the later mechanisms enforce those assumptions at measurable compute, communication, and coordination cost.

An attacker can steal database records, alter them, or disrupt access to them. Machine learning adds a learned representation through which those failures can propagate. An adversary may probe a model for memorized training examples or poison its training data to introduce a backdoor that activates on a chosen trigger. Protecting the stored records therefore does not fully protect the behavior learned from them or the information exposed through predictions.

Learned-information exposure follows from the difference between *transient processing* and *persistent learning*. Many applications process sensitive data without encoding it into learned parameters; machine learning systems can preserve patterns from training data in their model parameters. Sensitive information may then remain in a model after the original processing has ended and become accessible through outputs or systematic interrogation. Queries to a healthcare model may expose memorized patient information, while queries to a proprietary model may reveal enough behavior to train a substitute. The first threatens training-data privacy (section 13.5.2); the second threatens model confidentiality.

Operational platforms manage hundreds of models across distributed infrastructure, extending the attack surface across workers, devices, and administrative domains. In synchronous data-parallel training, one compromised worker contributes directly to the gradient that every replica applies, and the collective traffic carries that contribution across the training fabric. Its confidentiality depends on the fabric's isolation and encryption controls. Federated aggregation creates another observation boundary because model updates can reveal information even when raw records remain on devices.

Accelerators and caches shared among tenants introduce extraction and leakage risks that also exist on a single shared machine. Section 13.6 develops these fleet-internal boundaries. Scaling the system requires identifying which actors can observe or alter each exchange.

Architectural complexity compounds these challenges. A contemporary ML deployment spans data ingestion pipelines, distributed training infrastructure, model serving systems, and continuous monitoring frameworks, each introducing distinct vulnerabilities as figure 13.1 maps across the ML lifecycle. Continuous adaptation at edge nodes and federated coordination protocols add further entry points, each with its own trust assumptions.

↩ Multi-tenant platforms and the isolation they require are established in section 12.8.6.

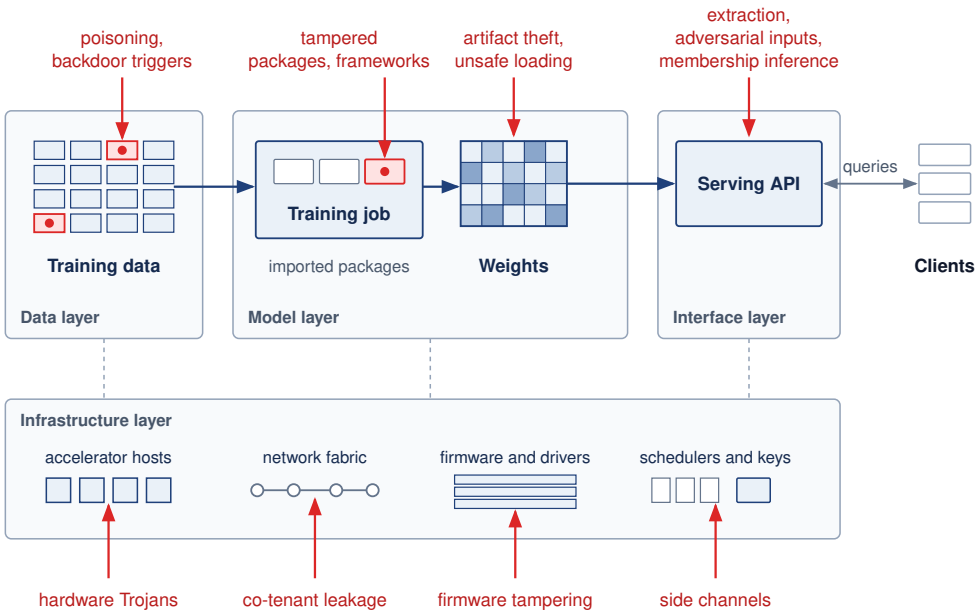

Figure 13.1: ML System Attack Surface: Red arrows mark where an adversary enters the system: the training data, the training job and its dependencies, the stored weights, the serving API, and the infrastructure beneath all three layers. Each entry point calls for a different control, so defense spans data collection, training and weights, the serving interface, and the provenance of firmware and hardware.

13.2 Security and Privacy Foundations

Defense starts by separating two concerns that often share the same vocabulary. Security asks how an adversary can steal, manipulate, or disable the fleet; privacy asks how sensitive information can leak or be inferred even when the system behaves as designed. Keeping those questions separate determines which controls belong at each layer: authentication and isolation around adversarial pathways, provenance and monitoring across the release chain, and privacy accounting wherever learned parameters carry traces of the data that produced them.

13.2.1 Security defined

The first half of that separation is security, which defends model parameters, training pipelines, deployment infrastructure, and inference interfaces against intentional adversarial behavior. An attacker who crafts an adversarial patch on a stop sign, steals model weights from an unauthenticated bucket, or floods a generative endpoint with cost-maximizing queries compromises the system's integrity, confidentiality, or availability.

Definition 13.1: Security

Security is the set of system properties (confidentiality, integrity, and availability) that protect an ML system's data, model weights, and inference pipeline from intentional adversarial actions, spanning both the infrastructure layer (network intrusion, credential theft) and the algorithmic layer (model extraction, prompt injection, adversarial examples).

1. **Significance**: Security failures operate on both surfaces simultaneously. At the infrastructure layer, theft of a large proprietary model represents direct IP loss. At the algorithmic layer, model extraction via black-box queries can recover decision boundaries or distill a functional approximation of a deployed model at a small fraction of its training cost (section 13.5.1 quantifies the query budgets), bypassing the investment and competitive moat. Either failure transfers the value of the original training investment to the attacker without the attacker paying its compute cost.
2. **Distinction**: Security specifies an intentional adversary and the boundaries that actor can cross. Robustness evaluates how model behavior changes under specified perturbations, including both accidental shifts and adversarial inputs; the two concerns overlap when the perturbation comes from an attacker.
3. **Common pitfall**: A frequent misconception is that traditional IT security (firewalls, access controls, encryption) adequately secures ML systems. ML introduces an algorithmic attack surface orthogonal to infrastructure: a properly authenticated API request containing an adversarial input or a prompt injection bypasses all network-layer defenses and manipulates model behavior through the model's own learned functions.

A facial recognition system deployed in public transit infrastructure, for example, may be targeted with adversarial inputs that cause it to misidentify individuals or fail entirely, representing a runtime security vulnerability that threatens both accuracy and system availability. The confidentiality-integrity-availability triad matters because each property breaks a different system contract. Confidentiality protects data and model access. Integrity protects the correctness of training and inference behavior. Availability protects the service from being disabled or degraded when it is needed.

13.2.2 Privacy defined

Whereas security addresses adversarial manipulation and unauthorized access, privacy limits the exposure and inference of sensitive information across training data, inference inputs, and model outputs. An authorized user querying a clinical classifier or a language model can extract a patient's sensitive diagnosis or verbatim personal record without ever breaching a firewall or stealing a password. Privacy protections must therefore account for unintended disclosure during ordinary operation. They must also account for active adversaries who choose queries, combine outputs with auxiliary information, or exploit repeated releases to infer sensitive facts.

Definition 13.2: Privacy

Privacy is the protection of sensitive information from unauthorized disclosure, inference, and misuse across the ML lifecycle.

1. **Significance**: It limits the exposure risk of training data and user inputs. Differentially private training introduces a utility-privacy trade-off: tighter privacy bounds generally require more noise or data, which can increase the training compute, the T_{compute} term of the Fleet Law (equation 1.3), required to reach a target utility.
2. **Distinction**: Unlike confidentiality (which focuses on access control), privacy in ML focuses on inference risks: the ability of a user, operator, or adversary to infer sensitive facts or reconstruct training samples from model outputs, weights, or repeated releases.
3. **Common pitfall**: A frequent misconception is that removing names (de-identification) is sufficient for privacy. In reality, neural networks are correlation engines that can inadver-

tently memorize and leak unique snippets of sensitive data through high-dimensional patterns.

Privacy failures are not limited to raw records or explicit identifiers. Aggregated behavioral traces can reveal sensitive routines when the surrounding geography makes individuals and sites easy to infer. Differential privacy (DP) bounds how much the distribution of released results can change when one protected contribution changes. Mechanisms can achieve this guarantee through calibrated randomization, such as adding noise to a query result. The protected contribution may be one record or all records from one person; the guarantee must specify which. Its utility cost depends on the task, data, mechanism, and training procedure.

Differential privacy provides one mechanism-level guarantee; it does not provide lifecycle privacy by itself. A system must still minimize collection, establish an appropriate basis for use, enforce access controls, limit retention, honor deletion obligations, and prevent raw inputs or sensitive outputs from leaking through logs and caches. These controls answer questions that a formal privacy bound does not: whether the data should have been collected, who may use them, how long derived artifacts remain, and which copies must be removed. Production privacy therefore composes formal accounting with data-governance controls across collection, training, serving, telemetry, and retirement.

Before the formal treatment of differential privacy, the chapter needs only one knob, the privacy budget ϵ. When one protected contribution changes, the probability of any released outcome changes by at most a factor of e^{ϵ}, apart from a small additive slack the formal definition allows, so smaller values give a tighter bound, and ϵ is not an accuracy percentage. Because additional data-dependent DP releases about the same data spend more of that budget, the budget belongs to the history of releases rather than to any one server. Using a fixed DP-trained model does not spend additional training privacy budget. Section 13.9 states the formal definition, its additive slack term, and the accounting that tracks spending across releases, once the defense architectures that surround it are in place.

13.2.3 Security versus privacy

Security and privacy intersect in controls such as encrypted storage, yet they differ in their objectives, threat models, and mitigation strategies. Encrypting a hospital's training bucket stops an external attacker from stealing files, whereas the model trained on that bucket can still reveal whether a specific patient was in the cohort through an ordinary, authorized prediction query. These security-privacy distinctions matter because the two domains optimize for different failure modes. Their distinct goals carry through to typical concerns, example attacks, the defenses practitioners must consider, and regulatory standing (table 13.1).

Table 13.1: Security-Privacy Distinctions: Machine learning systems require distinct approaches to security and privacy; security mitigates adversarial threats targeting system functionality, while privacy protects sensitive information from both intentional and unintentional exposure through data leakage or re-identification. This table clarifies how differing goals and threat models shape the specific concerns and mitigation strategies for each domain.

Aspect	Security	Privacy
Primary Goal	Prevent unauthorized access or disruption	Limit exposure of sensitive information
Threat Model	Adversarial actors (external or internal)	Passive exposure and active inference with auxiliary information
Typical Concerns	Model theft, poisoning, evasion attacks	Data leakage, re-identification, memorization
Example Attack	Adversarial inputs cause misclassification	Model inversion infers sensitive training-data properties
Representative Defenses	Access control, adversarial training	Differential privacy; federation with accounted privacy controls
Relevance to Regulation	Emphasized in cybersecurity standards	Central to data protection laws (for example, GDPR)

Systems Perspective 13.1: The privacy-utility trade-off

Security and privacy are deeply interrelated but not interchangeable. A secure system helps maintain privacy by restricting unauthorized access to models and data. Privacy-preserving designs can improve security by reducing the attack surface; minimizing the retention of sensitive data reduces the risk of exposure if a system is compromised.

However, their controls can impose competing requirements. Differential privacy bounds information disclosed by a release but may lower model utility. Encryption restricts who can inspect protected data, so an audit that needs those data requires an authorized inspection path and appropriate key custody. The design must preserve useful evidence while limiting what that evidence reveals.

Security and privacy therefore constrain the same system through different boundaries. Access control, cryptography, trusted execution, and differential privacy address different failure modes; their combination must protect model assets, limit data exposure, and retain the evidence needed to investigate failures. The breach histories make those boundaries concrete by tracing how aggregate releases, compromised components, and excessive access can expose information or alter system behavior.

13.3 Learning from Security Breaches

A public heatmap built for runners became a military-intelligence leak when sparse GPS traces outlined movement around sensitive sites. That privacy failure belongs beside three landmark security breaches because each case turns an abstract control objective into an operational constraint. Behavioral telemetry can identify people. Supply-chain compromise can alter physical systems. Weak isolation can expose safety-critical control paths. Default credentials can turn cheap devices into attack infrastructure. Although most of these incidents did not target ML systems directly, every failure mode they demonstrated has a direct analog in training pipelines, inference APIs, edge deployments, and privacy-preserving data products. The 2018 Strava heatmap shows how an aggregate data release can fail on privacy without a single security intrusion. The three security breaches that follow trace supply-chain, isolation, and endpoint failures.

War Story 13.1: The heatmap that revealed routines (2018)

Context: In late January 2018, public scrutiny of Strava's global activity heatmap, an aggregate built from users' uploaded GPS traces, revealed that it exposed activity at military bases (Russell 2018). The release showed patterns of movement without displaying a list of named participants.

Mechanism: Aggregation hides individuals only where many people contribute. In sparsely populated regions, the few bright tracks came largely from personnel exercising on and around bases, so the heatmap traced base perimeters and the routes people ran, even where a site's location was already known.

Impact: The map exposed an operational-security risk in consumer telemetry. An aggregate intended to help people find exercise routes also told an observer where personnel moved routinely, prompting reviews of how fitness devices were used near sensitive facilities.

Design response: A system releasing location aggregates must assess what sparse regions reveal, limit sensitive contributions, and make participation controls understandable. These are design requirements suggested by the incident, not evidence that a particular product change eliminated the risk.

Systems lesson: Privacy failures can occur without names, passwords, or model weights leaking. Aggregation alone does not establish privacy when the released pattern and its surrounding context reveal sensitive behavior.

13.3.1 Stuxnet and supply chain compromise

Discovered in 2010, the Stuxnet worm used four zero-day vulnerabilities (flaws unknown to the vendor and therefore unpatched) among its infection and propagation mechanisms, including spread through infected USB media. It targeted Siemens Step7 software and the programmable logic controllers (PLCs) it managed in industrial systems, including the air-gapped[1] environment at Iran's Natanz nuclear facility (Farwell and Rohozinski 2011). Altered control behavior accompanied by apparently normal telemetry demonstrated that operational health signals can conceal a compromise. An analogous ML failure occurs when poisoned training data or a backdoored model changes selected predictions while the inference server continues to meet its latency and availability targets. The common mechanism is a violation of behavioral integrity that ordinary service-health monitoring need not detect.

ML supply chains that ingest public packages, datasets, model weights, or firmware face four analogous vectors: compromised dependencies (malicious packages in PyPI and conda repositories), poisoned datasets on public platforms, backdoored model weights in model repositories, and tampered accelerator firmware. Each can introduce behavior through an artifact the pipeline accepts. Because a model's behavior depends on every contributor to its data, code, weights, and hardware, the model can be trusted only as far as a verifiable record of those contributors extends. That statement is the provenance bound on trust (principle 22), and it carries the Stuxnet lesson to every artifact an ML pipeline admits. Artifact signatures (cryptographic checks that show which key holder produced an artifact and that it has not changed since) and immutable provenance logs establish origin and integrity evidence; controlled dependency admission limits which artifacts enter the pipeline. Behavioral testing addresses a different question, because an authentic, unchanged artifact can still contain a backdoor. Each stage of the Stuxnet attack chain has a counterpart in the ML supply chain (figure 13.2).

[1] **Air-gapped**: Physical network separation does not eliminate transfer paths through removable media or authorized personnel. An isolated ML training cluster still needs a controlled process for admitting frameworks, datasets, and pretrained weights; each transfer must preserve the intended trust boundary.

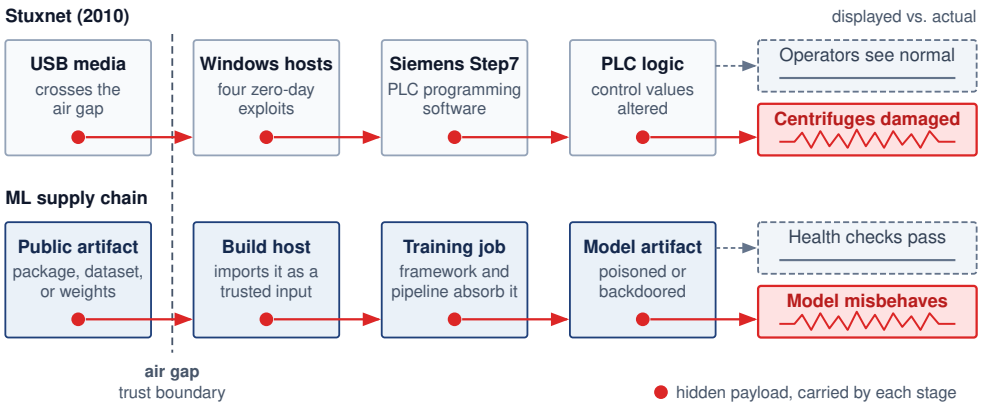

Figure 13.2: Stuxnet and the ML Supply Chain: The top row traces the Stuxnet chain from removable media through Windows hosts and Siemens Step7 software to altered PLC logic, which damaged centrifuges while operators saw normal readings. The bottom row places the ML analog under each stage: a compromised public artifact passes through trusted build and training steps into a model whose misbehavior monitoring does not reveal.

13.3.2 Insufficient isolation in the Jeep Cherokee hack

Security researchers remotely compromised a Jeep Cherokee's engine, transmission, and braking systems by exploiting a vulnerability in the vehicle's internet-connected Uconnect entertainment system, without physical access to the car (Miller and Valasek 2015; Miller 2019). Insufficient isolation allowed the compromised entertainment system to reach safety-critical controller area network bus controllers and manipulate vehicle functions.

The ML lesson concerns authority across a trust boundary. Perception or diagnostic models may need to communicate with actuator controllers, but compromise of an externally reachable inference

component must not grant unrestricted control authority. Segmentation, authenticated narrow gateways, sandboxed model execution, and minimal privileges constrain that path. Controllers must validate permitted commands and apply a system-specific safe response when inputs are invalid or communication fails; a network path makes the attack possible only if the controller accepts what arrives over it.

13.3.3 Weaponized endpoints and the Mirai botnet

In 2016, the Mirai botnet exploited weak credentials on Internet of Things (IoT) devices, including cameras, digital video recorders, and routers, to build infrastructure for distributed denial-of-service (DDoS) attacks. The longitudinal study of Mirai measured a peak of about 600,000 infections and documented attacks against services including Dyn, whose disruption affected other internet services (Antonakakis et al. 2017). The scaling lesson is that many individually limited devices can collectively threaten much larger services.

For ML edge deployments, a compromised endpoint may expose more than bandwidth. Its accessible records can include images or transcripts, and its enrolled federated-learning identity may let an attacker submit malicious updates. The defense must match those capabilities. Device-specific credentials limit shared-secret exposure, authenticated encrypted channels protect communication, and secure boot constrains which startup software is accepted. Hardware security modules (HSMs), tamper-resistant devices that hold keys and perform cryptographic operations without exposing those keys to host memory, can protect key custody where the platform supports them. Monitoring provides evidence of misuse, while update validation remains necessary because an authenticated device can still send harmful data.

Together, the Strava, Stuxnet, Jeep, and Mirai incidents expose failures at four different boundaries: an aggregate release, accepted components, paths to safety-critical control, and endpoint credentials. The engineering task is to identify the asset and attacker capability at each boundary before choosing a control. A threat model makes those assumptions explicit, so defenses can be evaluated against the failures they are intended to prevent and the resources they consume.

13.4 Systematic Threat Analysis and Risk Assessment

An authenticated request can still probe a model for memorized data, and an accepted training sample can still carry a backdoor. Network security and user authentication therefore need a threat model that includes learned behavior. Training data can be manipulated to embed backdoors, input perturbations can exploit learned decision boundaries, and systematic API queries can expose enough behavior to train a substitute. Each attack uses an interface the system needs for ordinary work.

Each of these vectors requires a formal **threat model**, written down as four fields: asset, boundary, adversary, and control. The asset defines what must remain confidential, intact, or available. The boundary defines where trust changes, such as between training data and model registry, user prompt and system instruction, or tenant workload and shared accelerator. The adversary field records the actor's capability (what it can access, observe, or modify) and its goal (what it seeks to compromise). The control identifies a permission, isolation mechanism, or observable signal that constrains the attack or supports a response. Without those fields, a risk matrix cannot explain why a particular defense deserves resources. Attack cost and expected benefit help assess feasibility, but the model must also cover actors whose objective is disruption rather than financial return.

For a federated training round, suppose the coordinator may inspect messages but follows the aggregation protocol, while a compromised enrolled client may submit an arbitrary update. The same update crosses two boundaries. Hiding individual contributions from the coordinator addresses confidentiality; checking or limiting a client contribution addresses integrity. Authenticating the client establishes who sent an update, not whether the update is benign. A defense comparison is meaningful only after stating which of these actors, capabilities, and possible collusions it covers.

13.4.1 Security threat modeling for ML systems

Systematic threat modeling identifies what must be protected and from whom. It is a structured approach to security analysis that characterizes the attack surface and guides defensive investments.

For machine learning systems, threat modeling must account for dependence on training data, statistical decision boundaries, and distributed deployment patterns.

13.4.1.1 Attack surface analysis

The attack surface of an ML system encompasses all points where an adversary can interact with or observe the system. It includes the lifecycle risks of ordinary software, such as compromised dependencies, as well as the data and learned state that determine model behavior. Decomposing this surface into data, model, interface, and infrastructure layers identifies where a control can constrain the attacker and what remains exposed beyond that control.

The training data pipeline lets an adversary influence behavior before the model exists. At collection endpoints, raw data enters through application logs, sensors, forms, APIs, or uploads. Training corpora then reside in object stores, warehouses, feature stores, or dataset registries. Preprocessing pipelines normalize, filter, augment, and join these inputs, while human annotation, weak-labeling rules, or model-assisted labeling supply the targets used for learning. Dataset manifests and lineage systems identify which snapshot enters a release. A compromised source can affect successive model versions if the pipeline continues to ingest its records, and a backdoor may persist when training reuses an affected checkpoint. Retraining alone therefore does not establish that the compromised contribution has been removed.

The model itself presents multiple attack surfaces spanning training infrastructure, model storage and versioning systems, model serialization and deserialization code, hyperparameter configuration management, and gradient computation and aggregation processes. Attacks at the model layer can compromise integrity by embedding trojans during training, extract intellectual property through parameter theft, or manipulate behavior through weight poisoning.

Deployed models expose additional attack surfaces through inference API endpoints and load balancers, authentication and authorization systems, input validation and preprocessing logic, output formatting and response generation, and monitoring and logging infrastructure. At this interface layer, input checks enforce the accepted request contract, query limits constrain collection, and output controls restrict the information returned. An adversarial input can satisfy the request contract, and a permitted response can still reveal useful model behavior, so these controls need evaluation against the specified attack.

The underlying computational infrastructure presents traditional attack surfaces amplified by ML-specific concerns, including accelerator firmware and drivers, container orchestration and scheduling systems, network communication between distributed training nodes, key management and secrets infrastructure, and supply chain for ML frameworks and dependencies. Infrastructure compromises can affect all systems running on shared resources, making this layer critical for multi-tenant ML platforms.

13.4.1.2 Threat vector classification

Threat-vector classification matters because the same model failure demands different defenses when the attacker has only API access, partial design knowledge, or full insider access. The three axes are access type, auxiliary knowledge, and attack timing.

Access type specifies how the adversary can inspect or interact with the model. Under black-box access, the available interface returns outputs for submitted inputs. It can support model extraction (Tramèr et al. 2016) and query- or transfer-based adversarial attacks (Nicolas Papernot, McDaniel, and Goodfellow 2016) without exposing the target's parameters. Gray-box access adds partial internal information, such as the architecture or training procedure. White-box access supplies the model's parameters and architecture, enabling direct gradient-based attacks (which follow the model's gradient with respect to its input to find a small change that alters the output) (Goodfellow et al. 2015) and parameter copying when the files are accessible. These labels describe model visibility; they do not grant permission to alter training data, replace weights, or modify the runtime. Those write capabilities must be stated separately. Each access level, which the literature also calls a knowledge level, has a typical attack it enables and a typical deployment where it arises (table 13.2).

Table 13.2: Adversarial Knowledge Spectrum: Model access defines the white-box, gray-box, and black-box distinction used here; training-data access is specified independently. An open-weight model can expose its parameters while keeping its training records private. The table categorizes these models by access level, typical attack methods, and common deployment scenarios, clarifying the practical challenges of securing machine learning systems.

Adversary knowledge level	Model access	Training data access	Attack example	Common scenario
White-box	Full access to architecture and parameters	Specified separately	Crafting adversarial examples using gradients	Insider threats, open-source model reuse
Gray-box	Partial access (e.g., architecture only)	Specified separately	Attacks based on surrogate model approximation	Known model family, unknown fine-tuning
Black-box	No internal access; only query-response view	Specified separately	Query-based surrogate model training and transfer attacks	Public APIs, model-as-a-service deployments

Auxiliary knowledge adds the background information available beyond the exposed interface. A black-box client may know the application domain, public architecture documentation, or examples resembling the training distribution. An insider may know the deployment configuration without being able to read every training record. Model visibility and auxiliary knowledge therefore overlap, but neither replaces an explicit inventory of what the adversary knows. Claims of full knowledge must identify whether they include training records, preprocessing, and defense settings as well as model parameters.

Attack timing determines where the defense can still intervene across the system lifecycle. Training-time attacks, such as data poisoning or backdoor injection, manipulate the learning process before the model artifact is produced. Deployment-time attacks compromise model distribution, serialization, or container installation before serving begins. Once in production, inference-time attacks exploit the deployed model through adversarial inputs or extraction queries. Finally, postdeployment attacks target online model updates, monitoring telemetry, or operational feedback loops after the model is already in service. The defense surface therefore moves in lockstep with the model lifecycle, from data controls to artifact integrity to serving-time containment.

Combining these dimensions defines a testable threat scenario. A public API adversary might query a deployed classifier, know its task and output format, and lack access to its parameters or training records. A training insider might inspect the model and alter a specified subset of examples before fitting begins. The second scenario permits poisoning because it grants a write capability, not merely because the attacker knows more about the model.

13.4.1.3 Defense strategy framework

Effective defense against the threat vectors identified by the threat model requires a layered strategy that addresses each attack surface while accounting for adversary capabilities. The defense framework operates on three principles: defense in depth, minimal attack surface, and fail-safe defaults.

No single defensive mechanism provides complete protection. The principle of **defense in depth** requires multiple independent layers such that compromising one layer does not grant full system access. For ML systems, this means combining data validation with model robustness techniques, access controls with output perturbation, and software defenses with hardware security mechanisms.

The principle of **minimal attack surface** treats every exposed interface, stored artifact, and network endpoint as a potential entry or observation point. An API should expose only the functionality and output precision its users need, because detailed scores provide additional observations of model behavior. Encryption at rest protects stored models and training corpora when the attacker cannot also obtain their keys. Separating training clusters from inference endpoints constrains lateral movement, while protected access logs support investigation of sensitive operations. Each control reduces a particular exposure; none removes the need to examine the paths that remain.

The principle of **fail-safe defaults** makes access depend on established permission rather than the absence of a detected threat. A request that cannot be authorized must not acquire privileges because a service failed. Behavioral anomalies require a separate response policy because they can

also arise from legitimate traffic or distribution shift. Depending on the application, the system may reject a request, defer to a verified fallback, suspend a release, or isolate an affected component while preserving essential service. The safe response follows from the consequences of both continued operation and interruption; it cannot be inferred from an anomaly score alone.

Each layer of the attack surface calls for its own defensive mechanisms and detection methods (table 13.3). Data defenses validate provenance, model defenses protect learned behavior and weights, API defenses govern query exposure, and infrastructure defenses anchor runtime trust.

Table 13.3: Defense Mapping by Attack Surface: Each layer of the ML system attack surface requires specific defensive mechanisms and detection methods. Effective security integrates protections across all layers while maintaining detection capabilities that can identify attacks that bypass preventive controls.

Attack surface	Primary threats	Defensive mechanisms	Detection methods
Data Layer	Poisoning, label manipulation, supply chain compromise	Input validation, provenance tracking, secure data pipelines	Statistical anomaly detection, data quality monitoring
Model Layer	Backdoor injection, parameter theft, trojan insertion	Secure training environments, encrypted model storage, access controls	Model behavior analysis, weight distribution monitoring
API/Interface Layer	Adversarial examples, model extraction, membership inference	Input sanitization, rate limiting, output perturbation, differential privacy	Query pattern analysis, confidence distribution monitoring
Infrastructure Layer	Side-channel attacks, firmware compromise, supply chain attacks	Trusted execution environments (TEEs), secure boot, network segmentation, dependency scanning	Hardware performance monitoring, integrity verification

Attack surfaces, threat vectors, and the defense mapping together enumerate what could go wrong and which control addresses each failure. They do not say which failure to address first, and no deployment can fund every control at once.

13.4.2 Threat prioritization framework

A team that spends months hardening against physical laser fault injection while leaving a public inference API without query budgets or loading untrusted pickle checkpoints (a Python file format that can run arbitrary code when loaded) from a public hub leaves its highest-probability attack paths wide open. Prioritization ranks the enumerated threats by how likely each is in the deployment and how much damage it would cause. Impact is not purely technical. Regulatory obligations raise the cost of some failures, as when an erasure request under the European Union's General Data Protection Regulation (GDPR)[2] obliges the operator to trace personal data through derived datasets, logs, and deployed artifacts. A prioritization matrix based on likelihood and impact turns the threat model into an allocation decision. Teams automate defenses for threats that are both likely and damaging, prepare for rare but severe failures, and avoid spending scarce engineering time on low-value controls.

Penalties make the regulatory part of impact concrete. In healthcare, where ML systems routinely handle patient data, a breach can violate the Health Insurance Portability and Accountability Act (HIPAA), with significant legal and ethical consequences. Similar regulatory risks apply globally. For personal data within its territorial and material scope, GDPR authorizes fines up to 2 percent of total worldwide annual turnover when an organization fails to implement appropriate security measures, and up to 4 percent for violations of its core processing principles or data-subject rights (European Parliament and Council of the European Union 2016).

The following placements illustrate a threat model rather than universal attack frequencies. Poisoning can fall in the high-likelihood/high-impact quadrant in a federated deployment that admits untrusted contributors and suffers severe consequences from manipulated predictions. Such a deployment needs controls on update integrity and a response path for suspected compromise. High-likelihood/low-impact extraction describes a different scenario in which public queries are accessible but behavioral copying creates limited competitive loss. Output restrictions and query budgets may then be proportionate controls, although their effectiveness still requires evaluation.

2 **GDPR (General Data Protection Regulation)**: Article 17 establishes a right to erasure under specified grounds and exceptions. For an ML pipeline, responding to an applicable request requires tracing personal data through source records, derived datasets, logs, and deployed artifacts. Whether a trained model requires remediation depends on the data it retains and the applicable obligation; deleting a source row does not establish that every derived copy or learned disclosure has been removed. Retraining and machine unlearning are possible technical responses whose effectiveness requires separate validation.

The lower-likelihood quadrants receive different treatment. Low-likelihood/high-impact hardware side-channel attacks describe a deployment in which obtaining the necessary observations is difficult but a successful attack could reveal valuable secrets. The required access depends on the channel; some attacks use a physical probe, while others exploit shared resources or remotely observable timing. In the public API scenario used for this matrix, membership inference attacks, which test whether a particular record was part of the training set (section 13.5.2.1), may sit in a lower-likelihood/lower-impact quadrant than direct integrity failures. That placement must be reassessed when the training data is sensitive or the exposed outputs make membership easier to infer. The matrix records those deployment assumptions rather than assigning a permanent rank to each attack family.

The framework guides resource allocation, as figure 13.3 summarizes. The example prioritizes accessible model and data threats before more specialized hardware attacks. A deployment with sensitive co-tenants or exposed edge hardware may require a different order. The allocation decision depends on the expected harm, the control's effectiveness, and its engineering cost, not on the attack category alone.

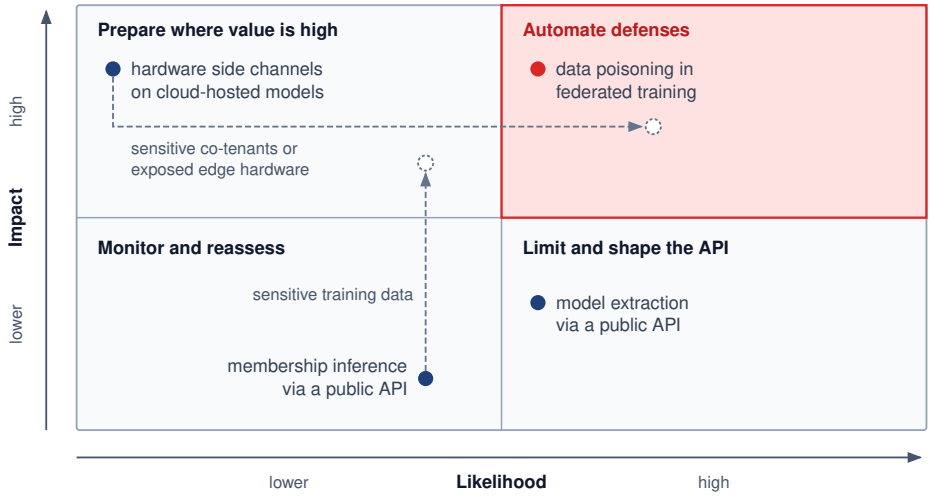

Figure 13.3: Threat Prioritization Matrix: Likelihood and impact place each example threat in a quadrant, and the quadrant sets the response. Likely, damaging threats such as federated data poisoning warrant automated defenses; rare but severe threats such as hardware side channels justify preparation in high-value deployments; likely but lower-impact extraction calls for API limits; and membership inference through a public API is monitored and reassessed when the training data is sensitive. Dashed arrows show how deployment assumptions move a threat to another quadrant.

Checkpoint 13.1: Knowledge check: Threat modeling

A team detects a high volume of API queries from a single IP address that are systematically exploring the decision boundary of a fraud detection model.

- ☐ **Lifecycle stage**: identify whether the suspicious activity occurs during training, serving, or infrastructure operation.
- ☐ **Attack surface**: identify the interface the attacker is exercising and infer what the query pattern reveals.
- ☐ **Defense mapping**: select the rate-limiting, monitoring, or output-control mechanisms that should activate first.

An unusual query pattern is evidence to investigate, not proof of extraction. The threat model connects that observation to an actor's capabilities and the asset at risk. Specific attack mechanisms then explain how access to data, artifacts, or predictions can compromise the model.

13.5 Model-Specific Attack Vectors

A stop sign with strategically placed black and white stickers can remain recognizable to a human while an evaluated image classifier labels it as a speed limit sign. The attacker changes the observed input rather than the classifier's code or stored weights. This adversarial example illustrates why the lifecycle entry point matters. Training data, model artifacts, and runtime inputs give an attacker different ways to influence the system and require controls at different boundaries.

The lifecycle entry point separates five objectives. Model theft targets confidentiality by obtaining artifacts or reproducing behavior from queries. Output leakage targets the training data, which the same queries can expose through the predictions the service is authorized to return. Poisoning targets training integrity by changing what the model learns. Adversarial examples target deployed behavior by changing the inputs a model receives. Availability abuse uses the same serving interface to exhaust the service's compute and cost budget rather than to change any prediction. Data poisoning therefore requires an intervention in the learning pipeline, whereas adversarial inputs and availability abuse call for controls on the deployed service.

Written as the four-field threat model of section 13.4, the objectives differ mainly in asset and adversary capability. Theft targets weights or behavior and needs either artifact access or query access, and output leakage targets the training data with nothing more than query access. Poisoning targets training integrity and needs write access to something the model learns from. Adversarial inputs target deployed decisions and need control over an input, while availability abuse targets the serving budget and needs nothing beyond an account. These objectives connect the attack surface analysis in section 13.4.1 to the controls available before training, during artifact handling, and at inference, while section 13.5.5 examines how token-stream architectures fuse instructions, untrusted content, and retrieved data across those objectives.

These entry points sit at different stages of the ML lifecycle (figure 13.4). Upstream, an attacker may inject samples or alter labels in a data source the pipeline accepts. A backdoor can then remain inactive on ordinary evaluation inputs and activate on a chosen trigger. Downstream, deployed files expose an artifact-access boundary, while prediction APIs expose a query boundary. An adversarial input crosses the latter without requiring the attacker to alter the training pipeline.

Controls must act at the relevant entry point and preserve their assumptions through later stages. Data validation limits what enters training, artifact controls authenticate what deployment loads, and serving controls constrain requests and returned information. None establishes behavioral correctness by itself; the threat model determines which attack each control must withstand.

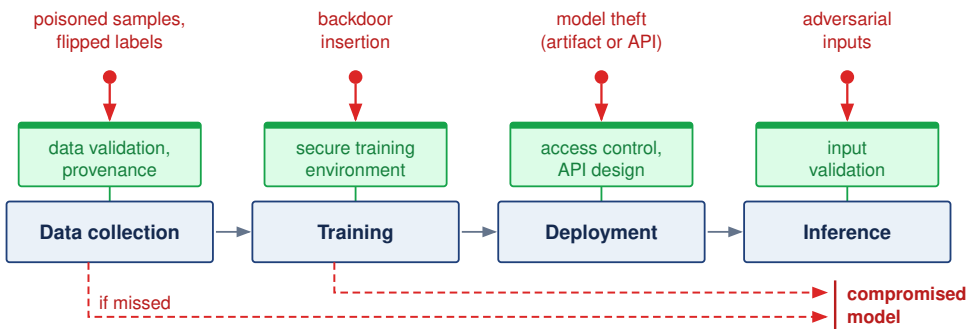

Figure 13.4: ML Lifecycle Threats: Poisoning, backdoor insertion, model theft, and adversarial inputs enter at different lifecycle stages (red arrows), and each stage has its own first control (green gates). Dashed lines show a training-time compromise that passes its gate persisting into every later stage, so a control placed at a different stage cannot intercept the threat before it takes effect.

13.5.1 Model theft

A competitor spending a few thousand dollars on public API queries to train a surrogate model, or an insider copying unencrypted weights from a bucket, can replicate a model that cost millions of dollars to train (Oliynyk et al. 2023). Model theft, the first of those five objectives, targets confidentiality by recovering a trained model's parameters, architecture, or output behavior, undermining its economic value or exposing private information encoded in its weights.

Such threats arise across a range of deployment settings, including public APIs, cloud-hosted services, on-device inference engines, and shared model repositories. Prediction interfaces can leak enough signal to recover model functionality (Tramèr et al. 2016; Oliynyk et al. 2023), while classifier outputs can reveal meaningful information about training sets or learned structure (Ateniese et al. 2015). Artifact handling introduces separate risks. Unsafe deserialization[3] can execute attacker-supplied code, while insufficient access controls can expose model files or deployment packages directly.

[3] **Model serialization**: Python's pickle format, common in PyTorch checkpoint workflows, can execute arbitrary code during deserialization, so loading an untrusted checkpoint is equivalent to running untrusted code. Formats such as SafeTensors store only tensor data and remove that execution path (Hugging Face 2026).

Model theft pursues two distinct objectives, and each requires different controls (figure 13.5). **Exact artifact theft** obtains the model file or deployment artifact through storage, registry, or device access. The attacker can reuse the recovered weights without repeating their training. **Approximate behavioral theft** uses API responses, such as labels, logits, embeddings, or confidence scores, to train a surrogate. It can exploit permitted queries without breaking authentication. The literature names this query-based branch in several ways. *Model extraction* and *model stealing* are synonyms for the attack, while *model distillation* names its usual training step, in which a surrogate learns from the victim's outputs. Its usual product is a behavioral surrogate, but the same query access can also recover selected internal parameters, a third outcome that the OpenAI results below demonstrate. Encryption, access control, and protected deployment paths restrict artifact access; signatures separately authenticate artifact integrity. Output limits, query budgets, auditing, and pricing constrain behavioral extraction, while watermarking can support attribution after copying.

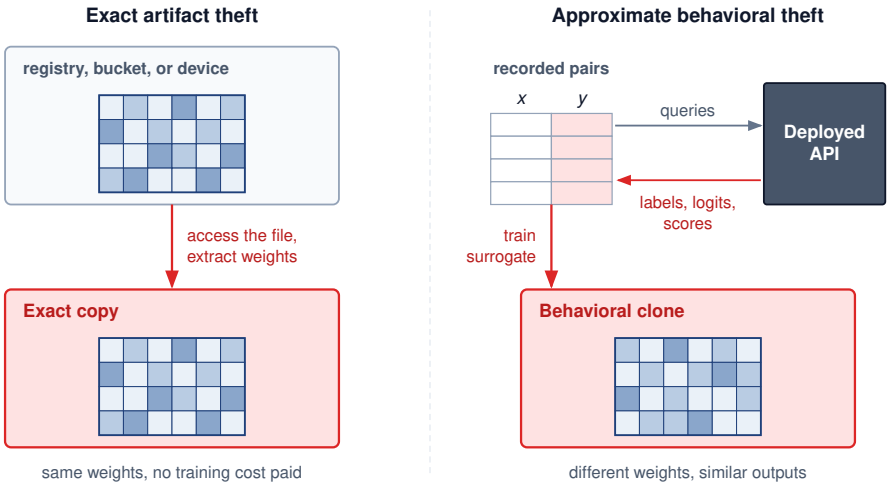

Figure 13.5: Model Theft Strategies: Exact artifact theft (left) obtains the model file, checkpoint, registry entry, or device artifact and reconstructs the original model. Approximate behavioral theft (right) uses black-box API queries, output signals, and surrogate training or distillation to clone the model's behavior without direct access to weights.

> Definition 13.3: Model extraction
>
> ***Model extraction*** is the class of attacks in which an adversary, with only black-box query access to a deployed model's prediction interface, recovers the model's parameters, decision boundaries, or behavior by analyzing observable outputs.

1. **Significance**: Successful extraction often costs orders of magnitude less than the original training. Published demonstrations have recovered logistic regression, neural-network, and decision-tree models from public machine-learning APIs with high fidelity and modest query cost (Tramèr et al. 2016). Later work has reconstructed components of large language models, including OpenAI deployments studied by Carlini et al. (2024), using only public query access. The BERT extraction example below illustrates the same dynamic on a transformer model, using only synthetic queries to train the copy.
2. **Distinction**: Extraction differs from direct theft of model files (which requires intrusion into the storage layer) and from model inversion (which uses model outputs to infer training-data properties or reconstruct representative inputs rather than the model itself). Its attack surface is the model's query interface, so every deployed ML service faces it by construction whenever the model is reachable.
3. **Common pitfall**: A frequent misconception is that rate limiting alone is a sufficient defense. An attacker willing to issue queries slowly, distribute requests across accounts, or accept a multi-month attack window may still extract useful behavior. Reducing output information, detecting systematic probing, and pricing queries can raise the cost or expose the attempt, but their combination does not guarantee prevention.

Model theft can also enable attacks on information encoded by the model. Extraction studies demonstrate recovery of useful behavior or selected internal parameters, although those outcomes do not imply complete replication of a large model or its training pipeline (Tramèr et al. 2016; Carlini et al. 2024). A competitor holding a stolen Archetype B (DLRM at Scale) recommendation model obtains its embedding tables, which store a learned vector for each value of each categorical feature. In a recommender those features commonly include user and item IDs, so the stolen tables hand the competitor a per-user vector shaped by that user's history, a natural starting point for inferring the user's attributes. Whether this exposes particular customer attributes depends on the model, accessible features, and auxiliary information. Model inversion attacks investigate that boundary by attempting to infer sensitive input properties from the learned model (section 13.5.2.2).

Example 13.1: The BERT model extraction

Scenario: Proprietary language models deployed behind public inference APIs return full probability vectors and confidence scores to client queries (Krishna et al. 2020).

Diagnosis: The attacker needed no real training data. Because both the victim and the copy were fine-tuned from the same publicly released pretrained BERT, queries made of random word sequences shaped by simple task-specific heuristics supplied enough labeled outputs to train a student. That student recovered much of the victim's task performance on tasks including natural language inference and question answering (Krishna et al. 2020). With a shared public starting point, the API responses had to supply only the task-specific fine-tuning signal. Retaining task accuracy is distinct from agreeing with the target on each input; neither implies recovery of its exact parameters.

Systems lesson: API endpoints expose intellectual property when outputs carry raw decision-boundary logits. Serving platforms combine rate-limiting, query auditing, and confidence-score truncation to make model extraction more costly and detectable.

13.5.1.1 Exact artifact theft

Exact artifact theft targets the model's stored structure and learned parameters. The exposed path may be an artifact-download API, an on-device inference package, or a model repository. This differs from querying a prediction API: the attacker obtains the artifact itself. Weak access control or unprotected packaging can expose that artifact without requiring full control of the underlying infrastructure.

The value of a stolen artifact depends on what it contains and what the attacker wants to reproduce. Learned parameters and architecture details can support inference when the necessary preprocessing

and runtime are available. Training hyperparameters, including learning rate schedules, batch sizes, and regularization settings, help reproduce or extend training, but are not required merely to execute an existing checkpoint. Side-channel attacks,[4] reverse engineering, or analysis of observable outputs may reveal additional structure when the artifact is incomplete. The threat model should therefore distinguish running a copied model from rebuilding the pipeline that produced it.

System designers must account for these risks by separating artifact security from interface security. Artifact controls secure model files, registry entries, and deployment packages; interface controls restrict APIs, reduce output precision, audit query patterns, and price or rate-limit access so behavioral extraction becomes harder to perform economically (Tramèr et al. 2016; Oliynyk et al. 2023).

[4] **ML side-channel attacks**: Timing or power observations can reveal properties of inference even when model files and query outputs are restricted. Recovering useful structure depends on the implementation, measurement access, and signal quality; physical proximity is a precondition for some channels, not a guarantee of success.

13.5.1.2 Insider theft of engineering assets

Insider access is often the shortest path to artifact theft, because an insider needs no exploit to reach files that the development environment already grants. The two cases below involve documents and source code rather than model weights, and that difference is their lesson. The protected asset extends beyond the serialized model to the code and data that produced it. Former Google engineer Anthony Levandowski was accused of taking proprietary Waymo materials before founding a competing startup (The New York Times 2017), a case about engineering intellectual property rather than demonstrated extraction of model weights. In March 2019, Tesla filed two trade-secret lawsuits involving former employees who left for autonomous-vehicle competitors (Korosec 2019). One complaint alleged that employees who joined Zoox copied proprietary warehousing, logistics, and inventory-control documents. A separate complaint alleged that former Autopilot engineer Guangzhi Cao copied Tesla Autopilot source-code repositories before joining XPeng/XMotors.

These allegations matter for ML systems because training code, feature pipelines, labeling workflows, deployment scripts, simulation assets, and operational playbooks can reveal enough about a system to accelerate a competitor's replication effort or expose downstream security weaknesses. API-level extraction defenses do not reach this boundary. Restricting prediction queries cannot prevent a developer from copying repositories that the developer can already access. Development access, artifact handling, and release provenance therefore belong in the same threat model as inference interfaces, with controls matched to the sensitivity of each asset.

13.5.1.3 Query-based extraction

Approximate behavioral theft uses observed input-output pairs to train a substitute without copying the target's parameters. This application of model distillation, or knockoff modeling, can reproduce useful functionality from public API responses (Orekondy et al. 2019).

Evaluation must distinguish task performance from agreement with the victim. A substitute can achieve similar benchmark accuracy while making different predictions; reproducing the victim's correct answers and mistakes measures behavioral fidelity. The service therefore has to limit accumulated observations while preserving the outputs legitimate users need.

Query access can also expose internal parameters rather than only supply examples for a surrogate. In their paper, Carlini et al. (2024) show how to reconstruct the final embedding projection matrix of OpenAI's `ada` and `babbage` models, and to recover the hidden dimensionality (and estimate the cost of full matrix recovery) for `gpt-3.5-turbo`, using only public API access. That final projection maps a hidden state vector to a much larger token vocabulary, so every output logit vector lies in a linear subspace whose rank equals the model's hidden dimension. Collecting logit vectors across enough varied prompts until new responses become linearly dependent on earlier ones reveals where the rank saturates and recovers the hidden dimensionality. For the smaller models, solving for a basis of that subspace replicates the weight matrix up to affine transformations, because the logits alone cannot distinguish the true hidden basis from any invertible transformation of it. This partial parameter extraction is distinct from both copying a complete artifact and training a behavioral substitute. It stops short of recovering the full model, yet it shows that an API's exposed scores and query controls can reveal internal structure that its operator intended to keep private.

The empirical results in table 13.4 demonstrate extraction of smaller models' output-projection parameters with root mean square errors on the order of 10^{-4}, while the larger GPT-3.5 rows show dimension recovery and estimated recovery costs rather than implemented weight-matrix extraction.

The production-interface attack used features such as top-k token log-probabilities together with logit-bias controls. The logit-bias control adds a caller-chosen offset to the logits of named tokens. Because a production API returns only the top-k log-probabilities rather than the full vocabulary vector, the caller can use that control to push selected tokens into the top-k set and read their relative logit values across prompts. The interaction between returned information and permitted query modifications therefore belongs in the threat model, even when neither feature exposes weights directly.

Table 13.4: Model Stealing Costs: Attackers can recover internal model information with publicly available APIs. The Cost column gives two figures: the dimension-recovery cost and the full weight-matrix extraction cost. For OpenAI's ada and babbage models, implemented output-projection weight extraction achieves low root-mean-square error with fewer than 4×10^6 queries at estimated costs of $1 to $12. The GPT-3.5 rows report dimension-recovery results and estimated full-recovery costs, not implemented weight-matrix extraction.

Model	Size (dimension extraction)	Number of queries	RMS (weight matrix extraction)	Cost
OpenAI ada	1024 ✓	$< 2 \times 10^6$	5×10^{-4}	$1 / $4
OpenAI babbage	2048 ✓	$< 4 \times 10^6$	7×10^{-4}	$2 / $12
OpenAI babbage-002	1536 ✓	$< 4 \times 10^6$	Not implemented	$2 / $12
OpenAI gpt-3.5-turbo-instruct	Not disclosed	$< 4 \times 10^7$	Not implemented	$200 / $2,000 (estimated)
OpenAI gpt-3.5-turbo-1106	Not disclosed	$< 4 \times 10^7$	Not implemented	$800 / $8,000 (estimated)

13.5.1.4 Defenses against model extraction

Unlike data poisoning and adversarial perturbations, whose algorithmic defenses are developed in chapter 14, and training-data leakage, whose formal privacy bounds are developed in section 13.9, model extraction is countered directly at the serving API boundary. Its quantitative defenses therefore belong beside the attack, while section 13.8.2 covers artifact custody, cryptographic packaging, and endpoint authentication infrastructure.

Model extraction defenses must account for what an adversary can accumulate through the API across four operational levers. Reducing output precision or returning fewer classes restricts the observations supplied by each response. Query quotas constrain collection per credential, and pricing increases its cost. Behavioral monitoring supplies evidence of systematic probing that can trigger investigation or enforcement. A daily rate limit only slows collection without bounding its lifetime total, and an attacker can combine observations from several accounts. The defense must therefore connect output design to cumulative budgets and the identities over which those budgets can actually be enforced.

Monitoring supplies evidence for adjusting those controls. High query volume is one possible signal, although legitimate batch jobs can produce it and an extractor can deliberately slow collection. A high-traffic, general-purpose prediction API might begin with an illustrative volume rule:

$$\text{alert}(\text{user}) = \begin{cases} 1 & \text{if } q_{\text{daily}} > 10{,}000 \text{ or } q_{\text{hourly}} > 2{,}000 \\ 0 & \text{otherwise} \end{cases}$$

where q_{daily} and q_{hourly} represent query counts. Such thresholds are only the first signal. They must be calibrated by service tier and by the shape of normal use, because batch customers and attackers can both generate high volume.

The next signal compares the inputs with expected use. Synthetic or systematically varied queries may reveal probing, but an extractor can also submit ordinary examples. A detector can compare the user's query distribution with a baseline chosen for the service and customer workload:

$$\mathcal{D}_{\text{KL}}(p_{\text{user}} \| p_{\text{expected}}) > \tau$$

where $\mathcal{D}_{\text{KL}}$ is Kullback-Leibler (KL) divergence, a measure of how far one probability distribution departs from another, and τ is an alert threshold. A practical detector needs a tractable representation

of inputs and enough observations to estimate those distributions. For a language model API, unusual token distributions or repetitive prompt templates may provide useful features. Regular inter-query intervals can add evidence of automation, but legitimate applications also automate requests and attackers can vary their timing.

An anomaly detector combines these signals into a score that can drive enforcement:

$$\text{score}_{\text{anomaly}}(u) = f_\theta(\text{qps}(u), \mathcal{D}_{\text{KL}}(u), \sigma_{\text{inter-query}}(u), \ldots)$$

where f_θ is a trained classifier and the features capture query volume, distribution, and temporal regularity. The score supports a response policy rather than establishing that theft occurred. A service can tighten quotas, request additional verification, or refer activity for review, with thresholds calibrated against legitimate workload variation. Free, paid, and enterprise tiers can carry different limits, but account-level enforcement constrains an adversary only to the extent that collection across credentials is also controlled.

A concrete quota ladder for the same general-purpose API makes that information cap visible. A free tier might allow 1,000 queries/day with a 10 queries/s burst, a basic paid tier might allow 100,000 queries/day with 100 queries/s, and an enterprise tier might allow 10M queries/day with custom burst limits and contractual monitoring. The exact numbers depend on product needs, but the ladder matters because extraction cost scales with accumulated responses.

An illustrative adaptive policy connects the anomaly score to a request limit. Higher scores reduce the quota, so false positives translate directly into lost capacity for legitimate workloads:

$$\text{limit}_{\text{effective}}(u) = \text{limit}_{\text{base}} \times \exp(-\gamma_{\text{rate}} \cdot \text{score}_{\text{anomaly}}(u))$$

where γ_{rate} controls sensitivity. With $\gamma_{\text{rate}} = 1$, a user with anomaly score 0.8 keeps about 45 percent of the base quota, a 55 percent rate reduction. The calculation specifies enforcement strength, not detection accuracy. Its security effect also depends on the response: full logits, hidden states, or detailed confidence scores can make each permitted query more useful to an extractor. Quota design therefore needs an explicit output contract.

Output shaping reduces the leakage per response while preserving the part of the answer legitimate users usually need. **Confidence rounding** returns coarser probabilities instead of full-precision logits or probabilities:

$$p_i^{\text{rnd}} = \text{round}(p_i, d_{\text{dec}})$$

where d_{dec} is the number of decimal places. For a 1000-class ImageNet classifier, rounding from 6 decimals to 2 decimals removes 4 decimal digits, about 13.3 bits per class. The count uses a decimal-bin model, which treats each decimal digit as selecting one of ten equally likely levels, so every dropped digit removes $\log_2 10 \approx 3.3$ bits. This counts representational precision, not measured information about the model. Rounding can also create score ties, so the service must evaluate the effect on ranking and downstream decisions.

The **top-k truncation** defense returns only the top-k predicted classes rather than the full probability distribution:

$$\text{output} = \{(c_i, p_i) : i \in \text{top-}k\}$$

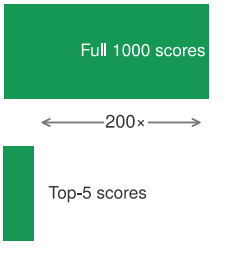

Top-k truncation exposes far fewer scores per query than the full distribution.

For $k = 5$ on a 1000-class problem, the API returns only those top scores instead of the full softmax vector, eliminating 99.5 percent of the distribution entries an extractor could otherwise observe per query. Legitimate users often need only the top few predictions, while an extractor benefits from the full distribution.

When rounded or truncated outputs still leak too much, calibrated noise can make repeated measurements less useful. Adding Gaussian noise to the score vector $\hat{\mathbf{y}}$ of a K-class model (uppercase K for the class count, distinct from the k of top-k) can push some scores below zero, so the noisy vector is projected back onto a valid probability distribution:

$$\tilde{\mathbf{p}} = \Pi_{\Delta^{K-1}}(\hat{\mathbf{y}} + \mathbf{z}), \qquad \mathbf{z} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I})$$

where $\Pi_{\Delta^{K-1}}$ returns the closest valid distribution to the noisy vector, one whose scores are all nonnegative and sum to one ($\tilde{p}_i \geq 0,\ \sum_i \tilde{p}_i = 1$). The set of such vectors is the probability simplex

Δ^{K-1}. The noise scale σ must balance extraction defense with utility preservation. Too little noise leaves the confidence surface measurable; too much noise changes decisions for legitimate users. This randomized response shaping is not, by itself, a differential privacy guarantee: such a claim also requires an adjacency relation, a sensitivity bound (the most one protected contribution can change the result), calibrated noise, and composition accounting.

Differential privacy addresses a different objective when a response is computed from sensitive records. It bounds the change in the released distribution under a specified change to the protected data. Section 13.9 develops the full (ϵ, δ) definition and the adjacent-dataset notation. Repeated interactive DP releases compose, but queries to a fixed DP-trained model are post-processing, computation on the released model that never reads the protected data again, and do not consume additional training privacy budget. Such queries may still enable model extraction, so output shaping and rate limits remain security controls.

The final lever is economics. Consider a simplified comparison in which extraction and independent training produce equally useful substitutes and query charges represent the extraction cost. If training from scratch costs C_{train} and extraction requires n_{queries} queries at price p_{query} per query, extraction is the cheaper option when:

$$n_{\text{queries}} \cdot p_{\text{query}} < C_{\text{train}}$$

Setting p_{query} so that $n_{\text{queries}} \cdot p_{\text{query}} > C_{\text{train}}$ makes extraction costlier than independent training in this simplified equal-quality comparison. It does not remove incentives based on unavailable training data, competitive value, or objectives other than acquiring a substitute cheaply. For example, if training a competitive large model costs $5M and extraction requires 100M queries, the break-even price is $0.05/query; pricing above that makes extraction more expensive than retraining under those assumptions. The comparison must include the attacker's alternative costs of compute, data, and engineering. In practice, pricing, monitoring, quotas, and output shaping work together. Restricted outputs and coarse confidence bins reduce exposed detail, while quotas constrain access through the identities they cover. Neither control alone bounds all information an actor can accumulate across accounts, alternate endpoints, or other sources.

Napkin Math 13.1: Protecting a production API

Problem: For a production ResNet-50 image-classification API serving 1M daily queries from 10,000 users, how can the operator make model extraction economically unattractive without degrading normal service?

Scenario: The model costs about $5K to train in this scenario. An extractor needs about 5M queries for 90 percent fidelity, or 5,000 queries/class for a 1,000-class classifier. Without defenses, an attacker abusing a free tier can send 25K queries/day for 200 days and pay nothing, making extraction economically favorable. The service is far smaller than the general-purpose API above, so its tiers and alert thresholds are scaled to its own traffic, and its training cost sits three orders of magnitude below the large model of the break-even example.

Mechanism: The defense stack works as a sequence:

1. **Cap collection**: The free tier allows 100 queries/day, the paid tier costs $50/month for 10K queries/day, and enterprise access requires authentication plus monitoring.
2. **Monitor probing**: The monitoring path watches for systematic measurement behavior, including more than 5K queries/day from one user or median inter-query gaps below 100 ms; the lightweight detector adds 2 ms per query.
3. **Shape outputs**: Output shaping rounds confidences to 2 decimal places, retaining about 6.7 bits from each rounded scalar probability, and returns only the top-5 classes, eliminating 99.5 percent of distribution entries.
4. **Price extraction**: At pay-as-you-go $0.001/query (above tier caps), 5M extraction queries cost $5K, matching the training cost in this scenario. An attacker limited to the $50/month tier (10K queries/day) would need 500 days and about $833 total, which is cheaper under these assumptions. A daily quota slows collection but does not cap its lifetime total.

5. **Cap cumulative collection**: A per-principal lifetime budget closes that route. If each principal's queries beyond an allowance well below 5M are billed at the pay-as-you-go $0.001/query, assembling the full extraction set costs close to $5K on any tier, the same as training the model.

Outcomes to measure: Whether the design works is decided by four measurements, namely the surrogate's accuracy, the accuracy legitimate users still receive, the latency the controls add, and the share of users flagged by mistake. The values below are constructed inputs that show the shape of that measurement, not results derived from the controls above. Suppose the surrogate's accuracy drops from 90 percent without defenses to 72 percent with defenses, an 18 percentage-point reduction, while legitimate top-5 accuracy remains 99.2 percent, close to the 99.3 percent baseline. Monitoring and output shaping add 2.5 ms on top of 100 ms inference, or 2.5 percent overhead, while flagging 0.3 percent of users for manual review.

Systems insight: Tier caps, monitoring, and output shaping slow an extractor; only the cumulative budget makes the full extraction set cost as much as retraining. Break-even is the design target rather than a weakness. At that price the query bill alone equals the cost of training an equivalent model. The extractor still pays for the surrogate's own training run and engineering on top of it while receiving a copy that the shaped outputs leave less faithful. Under the assumed equal-quality comparison, a lifetime price at or above break-even therefore removes the cost advantage of extraction over training an equivalent model independently. The residual route is to spread queries across many accounts, which account issuance controls and cross-account detection must bound, because a per-principal budget constrains an adversary only as far as principals map to real actors. The accuracy, latency, and false-positive figures are assumptions that the operator must measure before the design ships.

Pay-as-you-go pricing alone would charge the full training cost, but an extractor can stay inside the flat-fee tier, and tracking the cumulative bill against the queries collected shows what that choice is worth (figure 13.6). A daily quota stretches collection over 500 days but leaves the total near $833. A lifetime budget that bills every query beyond an allowance of 100K queries at the pay-as-you-go price instead makes the bill climb with the extraction set itself, reaching about $4,917 when the set is complete, close to the cost of training the model.

13.5.2 Leakage through authorized outputs

Whereas model theft copies the model, a second confidentiality failure needs no copy at all, because a model that executes correctly can disclose its training data through the outputs it is authorized to return. Persistent learning (section 13.1) carries sensitive information past the original processing, and the attacker reads nothing but permitted predictions, so storage encryption and caller authentication offer no protection. Three attacks exploit that leakage at increasing levels of detail. Membership inference learns a single bit indicating whether a given record was in the training set. Model inversion reconstructs properties of the training data, such as a representative input for a class or an identity. Training-data extraction recovers memorized records verbatim. All three need only query access or released weights, and all three grow more effective as the model fits its training data more closely than the population that data came from.

13.5.2.1 Membership inference

Membership inference tests whether a particular record was part of a model's training set (Shokri et al. 2017). The attack exploits a gap that training creates. The model has fitted its training members, so it usually assigns them lower loss and higher confidence than records it has never seen. The simplest test thresholds that signal and declares a record a member when the model's confidence on its true label exceeds a calibrated cutoff. A stronger test trains *shadow models* on data drawn from the same distribution with known membership. It learns from their outputs how member and nonmember responses differ, and it applies that classifier to the target model's outputs (Shokri et al. 2017). Both tests succeed in proportion to the gap between training and held-out behavior, so the risk grows with overfitting and tends to concentrate on small datasets and on rare records that the model had to fit individually. A control added for another threat can widen that gap,[5] so each one is

[5] **Robustness-privacy trade-off**: Adversarial training can increase susceptibility to membership inference in the settings studied by Song et al. (2019). The result means that a robustness intervention must be reevaluated against privacy invariants rather than assumed to improve every security property at once.

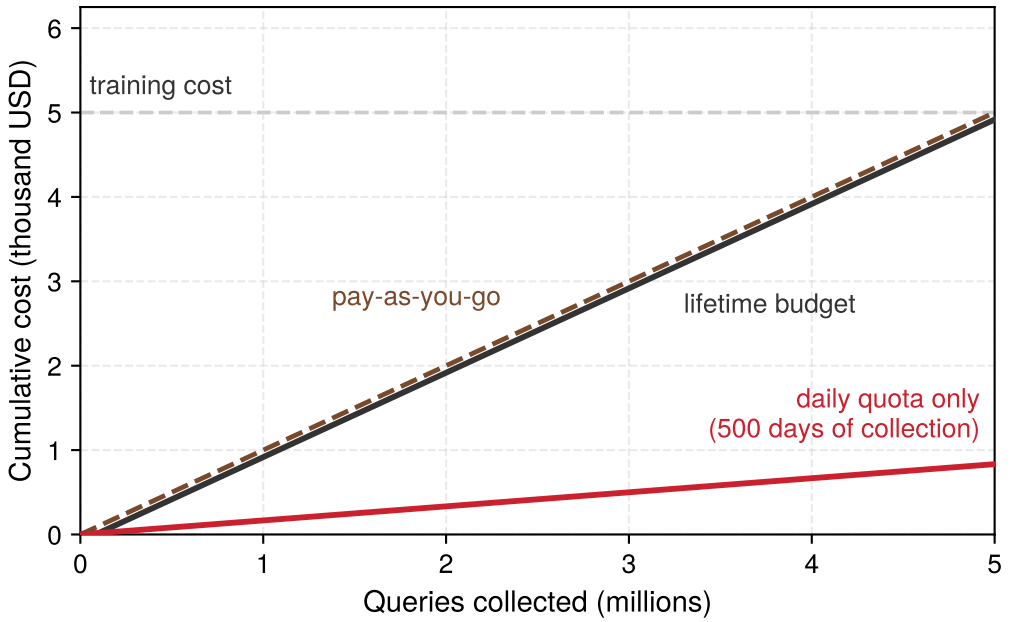

Figure 13.6: Extraction Cost under Three Billing Regimes: Cumulative query bill for the illustrative ResNet-50 API scenario as an extractor collects responses. Pay-as-you-go pricing and a lifetime per-principal budget both reach the training cost when the extraction set is complete; a daily quota on a flat-fee tier only slows collection, so the full set still costs a fraction of training. The scenario inputs are constructed, not measured.

rechecked against membership inference.

For threat modeling, the relevant variable is what the adversary observes: confidence scores, labels only, or individual client updates. Attack success depends on the target model, data, outputs, and adversary knowledge. Confirming that an individual's data was in a training set can still create privacy and compliance risk. In healthcare or legal prediction, membership itself can disclose sensitive information about a person, such as a diagnosis or a case. Whether that disclosure violates GDPR, HIPAA, or another rule depends on the data, parties, and deployment context (European Parliament and Council of the European Union 2016; U.S. Department of Health and Human Services 2003).

Membership inference is also the standard privacy audit, and the audit and differential privacy bound the same leakage from opposite sides. An audit runs the strongest available attack against the trained model and reports its success. It must state which attacker access it assumes, because some threat models permit only model queries while others grant parameters or auxiliary information. A successful attack proves leakage. A failed attack shows only that this attack, under these assumptions, found none, so an empirical audit is a lower bound on what the model reveals.

Differential privacy (section 13.9) supplies the upper bound. A membership test is itself a computation on the released model, so the guarantee applies to its verdict directly and limits how much more often any test can flag a protected record when that record was in training than when it was absent. Under a guarantee with budget ϵ and no additive slack, the rate at which a test flags true members can exceed the rate at which it falsely flags nonmembers by at most a factor of e^{ϵ}, which places a ceiling over every possible attack. The distance between the measured lower bound and the accounted upper bound is the uncertainty that a privacy claim carries, the band between the audit curve and the ceiling in figure 13.7.

13.5.2.2 Model inversion

Model inversion infers sensitive properties of the training data, or a representative input for a class or identity, from the learned model (Fredrikson et al. 2015). It can use either query access or exposed model internals. Confidence scores, for example, can guide a search for inputs that the

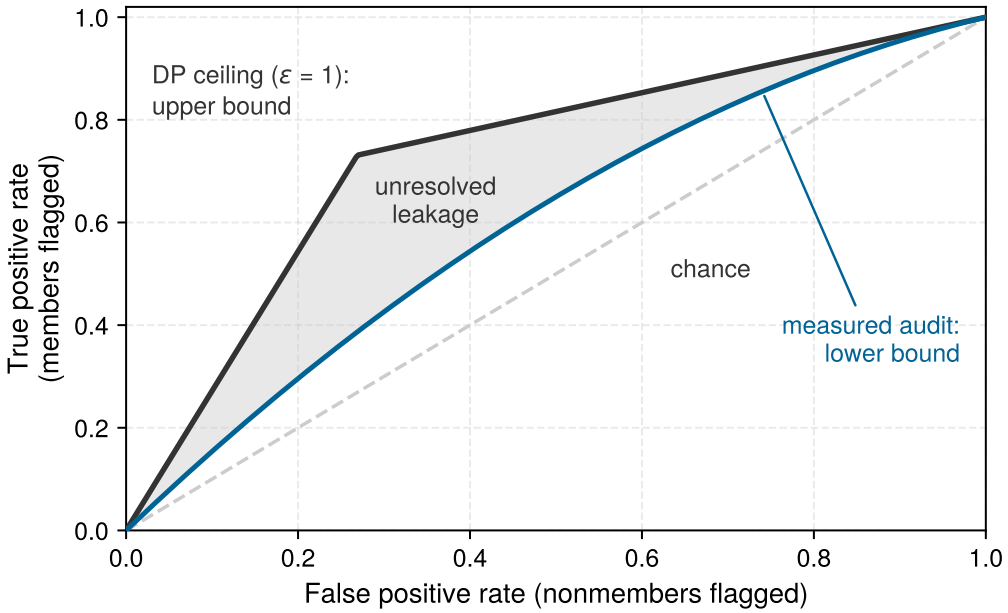

Figure 13.7: Membership Audit and the Privacy Ceiling: Each point on the plane is a membership test, read by how often it flags true members against how often it falsely flags nonmembers. A differential privacy guarantee with no additive slack at the illustrative budget $\epsilon = 1$ confines every possible test below the dark ceiling, an upper bound on leakage. An audit attack traces a curve that shows leakage it actually found, a lower bound. The shaded band between them is leakage that neither has ruled in or out. The audit curve is illustrative.

model associates with a target identity or attribute. A reconstructed face may reveal recognizable features without reproducing an exact training photograph; the privacy claim must distinguish those outcomes. Gradient inversion applies the same idea to the updates that a federated client sends (section 13.8.5.2).

13.5.2.3 Memorization and training-data extraction

Memorization is the strongest form of the leak, in which a model retains enough of a specific training record to reproduce it. It is not confined to language models. Under one published extraction protocol, a diffusion model reproduced a training-set photograph almost exactly (Carlini et al. 2023); figure 13.8 pairs the original with the reconstruction. Evidence of extraction concerns that specific training instance; a similar-looking output of the same subject would not establish memorization. The exposure matters most for assistants trained or adapted on user logs, where the memorized strings are the users' own messages.

13.5.2.4 GPT-2 data extraction case study

The GPT-2 extraction case shows why storage access control is insufficient once a model has learned sensitive content. In 2020, researchers demonstrated that query access to OpenAI's GPT-2 model could reveal verbatim training examples (Carlini et al. 2021). The attack targeted information exposed through generation rather than direct access to the stored training dataset. The researchers generated hundreds of thousands of samples, some unconditioned and some seeded with short snippets of internet text. They then ranked the samples by how much more likely GPT-2 found them than a reference model or compression baseline did. This allowed them to extract personal information such as names, phone numbers, and email addresses, along with source code, unique identifiers, and other verbatim snippets that the model had memorized from its training data.

The technical approach exploited GPT-2's memorization of rare or repeated text sequences through three stages. Sampling came first. The researchers generated 600,000 samples using top-k sampling, sampling with a temperature that decays over the first generated tokens, and sampling conditioned on short prefixes drawn from internet text. Ranking then applied membership-inference metrics

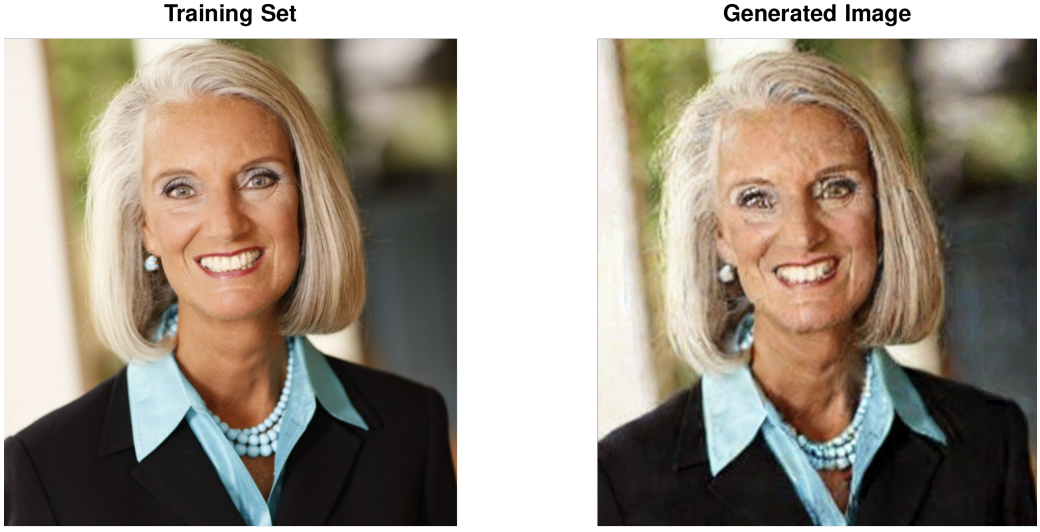

Figure 13.8: Diffusion Model Memorization: A training-set photograph (left) and a diffusion model's near-identical generation (right) under the extraction protocol of Carlini et al. (2023). Memorization is not confined to language models, so a model trained on sensitive images can return one of them through an authorized output.

that compare GPT-2's perplexity on a sample (a measure of how unlikely the model finds the text) against a reference, namely the sample's zlib compression entropy, the perplexity a smaller GPT-2 model assigns it, and GPT-2's perplexity on the lowercased text. A sample that GPT-2 finds far more likely than the reference suggests memorization rather than fluent generalization. Manual verification closed the pipeline. The researchers inspected 1,800 top-ranked candidates, checking them first by internet search and then against the training data with OpenAI's cooperation, and confirmed 604 of them as unique memorized training examples.

The attack challenged the assumption that training data would be "forgotten" after optimization. Some training examples remained recoverable even though the model was designed to generate text rather than retrieve records. Across the GPT-2 family studied (124M to 1.5B parameters), larger models exposed more memorized training data under the evaluation. This comparison motivates testing memorization as models scale; it does not establish a universal relationship between parameter count and privacy loss across different datasets and training procedures. For Archetype A (GPT-4/Llama-3), the deployment must therefore assess its actual corpus, training procedure, and exposed outputs rather than infer leakage from parameter count alone.

Filtering, deduplication, and manual review can reduce exposure, and output filters can screen known sensitive patterns, but empirical screening does not prove that no sensitive string remains recoverable. Differential privacy addresses a different requirement by bounding how much a protected contribution changes the distribution of released results. It does not prohibit every reproduction of text that appears in training, especially information shared across many records. Its contribution bound and the utility effects of filtering or private training must therefore be evaluated separately.

The result demonstrates why training-data controls need layered mitigations rather than storage access control alone (Carlini et al. 2021). Those layers can include filtering of personally identifiable information (PII), deduplication, extraction-style memorization audits, membership-inference testing, and differential privacy when a formal bound is required. Machine unlearning and policy controls address related lifecycle questions, but their guarantees require separate evidence and validation.

13.5.3 Data poisoning

While model theft and output leakage compromise confidentiality through read access to a finished model, data poisoning compromises training integrity through write access to something the model will learn from, arriving earlier in the lifecycle than any serving control can act. Injecting a few hundred mislabeled or trigger-stamped examples into a web-scraped corpus or user-feedback stream can plant a targeted backdoor while overall validation accuracy looks completely normal (Biggio et al. 2012).

Definition 13.4: Data poisoning

Data poisoning is the corruption of training data to compromise model behavior at inference time, either by injecting malicious samples or modifying existing labels.

1. **Significance**: It undermines the foundational assumption of data integrity. Even a small fraction of poisoned samples can create backdoors or systematic biases that remain latent until triggered by specific inputs during serving.
2. **Distinction**: Unlike inference-time evasion attacks against a fixed model, data poisoning operates through data collection or training, contaminating the model's learned mapping from the source.
3. **Common pitfall**: A frequent misconception is that poisoning can be "fixed" by more data. In reality, poisoning often exploits the aggregation property of training: adding more clean data may not "wash out" a carefully targeted backdoor that uses a unique trigger.

Poisoned data may cause a model to make incorrect predictions, degrade its generalization ability, or embed failure modes that remain dormant until triggered after deployment. The exposure is greatest in applications that retrain on data collected from external sources, including user interactions, crowdsourced annotations, and online scraping, because an attacker can then contribute training data without any access to the training pipeline itself.

The access axes of section 13.4.1.2 apply with one change of emphasis. Model visibility still determines how precisely a poison can be crafted. A white-box attacker can optimize examples against the actual architecture and training procedure, while a black-box attacker relies on surrogates and open submission channels. What makes poisoning possible at all, however, is the write capability that the threat-vector scenarios separated from visibility. That capability can sit at collection, preprocessing, labeling, or storage, so the poisoning surface spans the whole data path. Its priority relative to other threats follows the deployment context analyzed in section 13.4.2.

A poisoning attempt spans admission, learning, and exploitation. The attacker first places manipulated examples where the training pipeline may accept them. Training then determines whether those examples induce the intended behavior; merely inserting a record does not guarantee success. If the altered behavior survives evaluation and deployment, the adversary may exploit it to cause selected errors or degrade service. The attack must therefore be evaluated through the actual learning procedure and release gates.

Data poisoning can be viewed as a bilevel optimization problem, because the attacker's choice affects a model through the defender's training procedure. In the outer problem, the attacker selects poisoning data $\mathcal{S}_{\text{poison}}$ to increase loss on a validation or target dataset $\mathcal{S}_{\text{test}}$. The inner problem trains the model on the resulting data. Let $\mathcal{S}$ represent the original training data. Subject to the attacker's permitted modifications and poisoning budget, an untargeted objective is:

$$\max_{\mathcal{S}_{\text{poison}}} \mathcal{L}(f_{\mathcal{S}\cup\mathcal{S}_{\text{poison}}}, \mathcal{S}_{\text{test}}) \tag{13.1}$$

where $f_{\mathcal{S}\cup\mathcal{S}_{\text{poison}}}$ represents the model trained on the combined dataset of original and poisoned data. For targeted attacks, this objective can be refined to focus on specific inputs x_t and target labels y_t:

$$\min_{\mathcal{S}_{\text{poison}}} \mathcal{L}(f_{\mathcal{S}\cup\mathcal{S}_{\text{poison}}}, x_t, y_t) \tag{13.2}$$

In the training system, the nested objectives become two coupled loops (figure 13.9). The outer loop is the attacker's proposal step, which optimizes the attack objective, either raising loss on the test set (equation 13.1) or, for a targeted attack, lowering loss on the chosen label y_t (equation 13.2). The inner loop is the defender's training process, which minimizes loss on the combined data. The attacker evaluates or approximates this inner effect on training, and uses the resulting target loss to refine the proposal. An attacker without access to the defender's training process must instead rely on a surrogate or other available feedback.

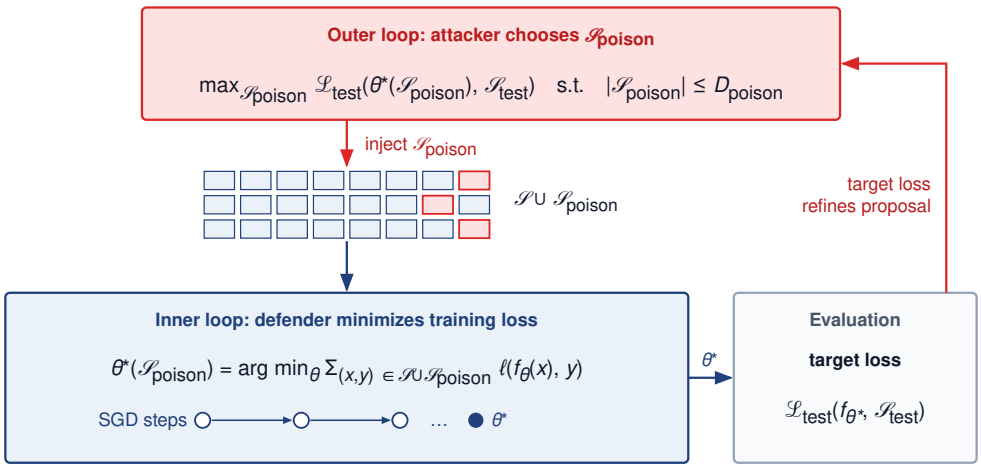

Figure 13.9: Data Poisoning as Bilevel Optimization: The attacker (outer loop) chooses poisoning data $\mathcal{S}_{\text{poison}}$ to raise the trained model's loss on a test set, the untargeted form drawn here, or, for a targeted attack, to lower its loss on the attacker's chosen label. The system (inner loop) attempts to minimize its training loss on the combined dataset $\mathcal{S} \cup \mathcal{S}_{\text{poison}}$. The nested attack objective can require repeated training evaluations or approximations. It does not imply that every poisoning defense must simulate full training.

This formulation captures the adversary's goal of introducing carefully crafted data points to manipulate the model's decision boundaries. Consider an email spam filter that retrains on messages users report as legitimate. An attacker who controls a set of accounts can report a batch of messages written in the style of an upcoming campaign as "not spam." The goal is to shift the filter's decision boundary so that the campaign's future messages reach inboxes. In this case, the poisoning data $\mathcal{S}_{\text{poison}}$ consists of mislabeled campaign-style messages, and the attacker's objective is to have the campaign messages x_t classified as legitimate mail y_t, following the targeted attack formulation in equation 13.2. If the mislabeled messages influence training enough to induce that behavior, aggregate performance on other traffic may still appear normal. The contrast with the stop-sign stickers that opened the discussion of model-specific attacks is the lifecycle stage. Poisoning changes the model before deployment, whereas the stickers change only the input that an unmodified deployed model receives.

Data poisoning attacks differ in their objective and scope (Biggio et al. 2012). Availability attacks seek broad degradation of useful performance. Targeted attacks seek errors on selected inputs or classes and may leave aggregate accuracy largely unchanged (Shafahi et al. 2018). The targeted objective of equation 13.2 rewards errors only on the attacker's chosen inputs, while the inner training loop still fits the combined data, which remain mostly clean, so the objective itself does not push aggregate accuracy down. Backdoor attacks associate an attacker-chosen trigger with unwanted behavior; the trigger can be visible, as in a sticker, even when its significance is hidden (Gu et al. 2017). Subpopulation attacks instead target a group defined by shared features (Jagielski et al. 2021). These categories can overlap. Their objectives guide which evaluations are needed beyond aggregate accuracy, including targeted, trigger-specific, and subgroup tests.

The objective classes say what an attacker wants; the attack *mechanism* depends on the attacker's access to the system and knowledge of the data pipeline. Every mechanism shares the shape traced

in figure 13.10, in which a poisoned record enters the training pipeline alongside clean data and attempts to alter the resulting model's behavior. The most direct mechanism is label modification, where an attacker selects a subset of training samples and alters their labels, flipping $y = 1$ to $y = 0$ or reassigning categories in multi-class settings. Even small-scale label corruption can shift decision boundaries. At web scale the same mechanism mislabels multimodal pretraining pairs, as in the dirty-label sample of figure 13.11, which pairs cat images with misleading "dog" captions to induce false associations.

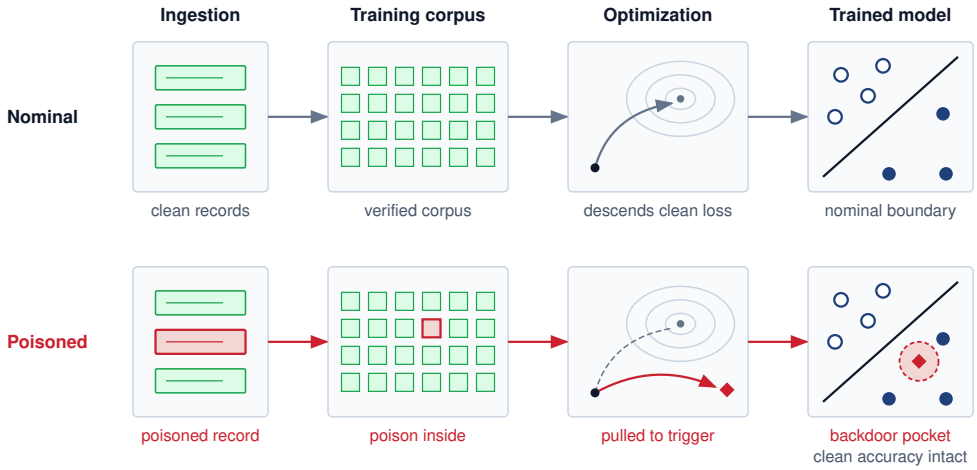

Figure 13.10: Data Poisoning Impact: A nominal pipeline (top) and a poisoned pipeline (bottom) pass through the same four stages. In the illustrated successful attack, a poisoned record admitted at ingestion contaminates the training corpus, pulls the optimizer's gradients toward the attacker's trigger, and yields a backdoored model whose clean accuracy can remain intact, so aggregate validation alone can miss it.

Figure 13.11: Dirty-Label Data Poisoning: Mismatched image-text pairs represent a dirty-label poisoning attack where cat images are paired with misleading 'dog' captions in training data. Incorporating these corrupted pairs can cause models to learn incorrect multimodal associations. Source: (Shan et al. 2023).

Beyond label modification, the attacker can keep the label intact and corrupt the features instead, the strategy known as **clean-label poisoning**. In the bilevel formulation, only the attacker's permitted modifications change, from labels to features. Because every label is correct, label audits cannot flag these records. Imperceptible image perturbations, subtle shifts in structured fields, and fixed trigger patterns aim for the same outcome. In each case the training example still appears legitimate, but it bends the learned decision boundary toward a future failure. Generative methods and data-synthesis tools raise the same risk at larger scale because they can create natural-looking examples whose only purpose is to distort what the model learns.

Clean-label poisoning is not always malicious in intent; researchers have begun exploring it as a defensive tool, particularly for protecting creative work from unauthorized use by generative AI models. Nightshade, developed by researchers at the University of Chicago, aims to deter

unauthorized training by poisoning images that a scraper may collect (Shan et al. 2023). It does not block scraping itself. Its perturbations preserve the apparent image content while changing how a model learns from it; the reported experiments show targeted concept corruption when the modified images enter the tested training pipelines.

The dose-response matrix in figure 13.12 charts this progressive corruption across four poison doses (0 to 300 samples) and eight prompt concepts. As the dose increases, Stable Diffusion XL outputs morph away from source concepts (generating cows for "car" or cats for "dog") under the tested training conditions. These outputs demonstrate changed behavior without, by themselves, identifying the internal representation that caused it. The experiment demonstrates concept poisoning, in which poisoned samples distort a model's semantic associations.

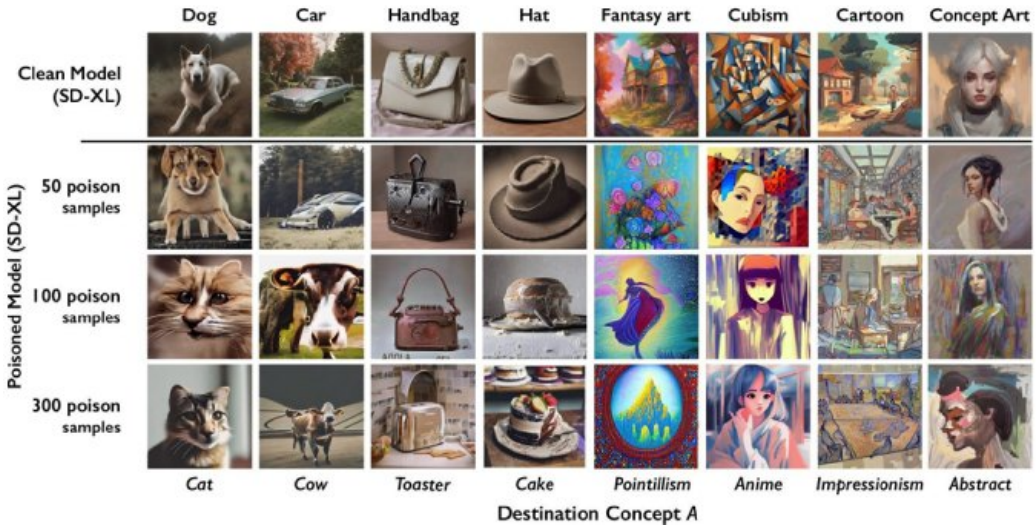

Figure 13.12: Concept Poisoning Dose-Response: Stable Diffusion XL outputs across four poison doses (clean, 50, 100, 300 samples) and eight source prompts (dog, car, handbag, hat, fantasy art, cubism, cartoon, concept art). As the dose grows, outputs morph from the intended source concept toward the attacker's destination concept, demonstrating Nightshade-style concept poisoning of a pretrained model. Source: (Shan et al. 2023).

The Nightshade experiments also report effects on concepts related to the poisoned target. This spillover extends the observed attack beyond the exact target prompt, but its reach depends on the model, poison construction, and training procedure. The same technique used to protect artistic content could also be repurposed to sabotage legitimate training pipelines, highlighting the dual-use dilemma[6] at the heart of machine learning security.

Whether crafted examples become poisoning depends on the data boundary. Web scraping, social media feeds, crowdsourced annotation, and untrusted user submissions all allow poisoned records to enter upstream and pass through weak cleaning checks in a "trusted" form. Physical systems add a second path. A sticker on a road sign attacks a car at inference time, but photographs of that sign harvested into a training corpus poison the model only when they are crafted or labeled to serve the attacker. Correctly labeled examples can instead help the model resist the sticker. Online learning systems tighten the feedback loop further, allowing an attacker to introduce small increments of malicious data until model behavior drifts without a single obvious anomaly. In collaborative training the same boundary sits at aggregation, where one malicious client's update reaches every participant (figure 13.13). Section 13.6.1 examines the same exposure among the workers of a synchronous data-parallel job.

Mitigating poisoning requires controls from data admission through model evaluation, and section 14.7 develops those defenses and the evidence for them. Because targeted and backdoor attacks can preserve aggregate validation accuracy, clean evaluation cannot certify integrity. The provenance bound on trust (principle 22) makes source verification and lineage necessary evidence about what entered training, but a verified source can still supply harmful records. Validation and anomaly detection examine those records, while robust training can limit their influence under

6 **Dual-use dilemma**: The structural tension that defensive ML capabilities (adversarial training, data poisoning tools, red-teaming frameworks) are simultaneously offensive capabilities. This creates an arms race where defensive research publications become attack playbooks. For ML systems engineering, the consequence is architectural: robustness mechanisms must assume attacker knowledge of the defense, ruling out security-through-obscurity and motivating formal guarantees where feasible and adaptive empirical evaluation within the remaining threat model.

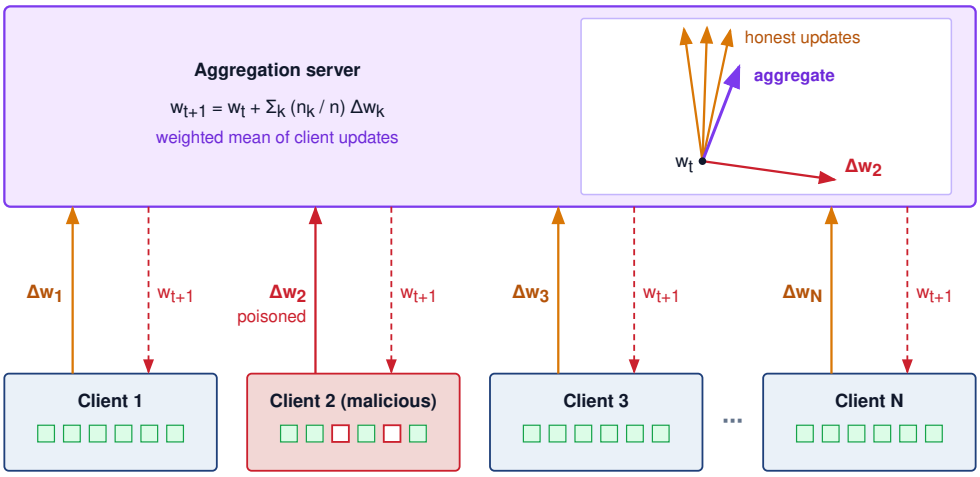

Figure 13.13: Collaborative Poisoning via Aggregation: In a collaborative learning setting, multiple clients each train on local data and contribute model updates to a central server that aggregates them into a shared global model. A single malicious participant can inject poisoned updates that the server folds into the aggregated model, so the corruption spreads to every client without any one obvious anomaly. This attack surface differs from adversarial examples because it targets the model during *training* rather than at *inference*.

a specified contamination model. Release evaluation must then test the targeted behaviors that ordinary accuracy metrics can miss.

↪ Poisoning defenses, from data sanitization to robust training, are developed in section 14.7.

13.5.4 Adversarial attacks

An attacker who can modify a sensor input or submit an API request can target the deployed model without changing its training data. The security question is which inputs the actor can control and which service decisions those inputs can influence. Behavioral robustness asks how much the model's predictions degrade under that specified manipulation. Here the threat model establishes the allowed attack; section 14.6 develops the evaluation and training methods needed to measure and improve resistance.

Adversarial inputs exploit the deployed model's decision behavior under attacker-chosen modifications. Some attacks use small, hard-to-perceive perturbations; others use visible patches or changes allowed by a physical or semantic constraint. The threat model must state those constraints before attack success or defense quality can be measured. This shifts attention from the integrity of upstream training data to the inputs and decisions exposed during inference.

These attacks create real-world risks in domains such as autonomous driving, biometric authentication, and content moderation. Small, visually hard-to-distinguish perturbations can cause image classifiers to make high-confidence errors (Szegedy et al. 2013; Goodfellow et al. 2015). Physical-world work showed that printed adversarial images can survive camera capture and remain misclassified by an ImageNet classifier (Kurakin et al. 2017). Road-sign attacks then showed that small stickers or perturbations can cause signs to be misclassified under realistic viewing conditions (Eykholt et al. 2018). A deployed defense must be evaluated against attacks adapted to its mechanisms. Changing the attacker model or the deployment can invalidate earlier evidence, so each change requires renewed testing.

The Perspective API case makes the inference-time distinction concrete. Researchers studying Google's toxicity detection API in 2017[7] showed that subtle misspellings, character substitutions, and added punctuation could cause explicitly toxic phrases to receive low toxicity scores (Hosseini et al. 2017). The training pipeline had not been poisoned; the deployed decision surface was fragile to small input perturbations.

The shared abstraction is a search over permitted input changes for an outcome the attacker wants. A norm bound constrains numerical distance, while a physical constraint may limit sticker placement

7 **Perspective API**: Google's toxicity model has been reported in use across many platforms, including *The New York Times* and Wikipedia. A shared classifier can create correlated exposure to an evasion technique, but success at each site also depends on its model version, preprocessing, thresholds, and additional moderation controls.

or camera conditions. Neither automatically establishes that a modification is imperceptible or harmless to a human observer.

The access types of section 13.4.1.2 determine how that search proceeds. White-box access lets the attacker compute input gradients directly, while black-box and gray-box attackers must estimate them from queries, partial architecture knowledge, or a surrogate model. Common attack strategies include surrogate model construction, transfer attacks exploiting adversarial transferability[8] (Nicolas Papernot, McDaniel, and Goodfellow 2016), and generative adversarial network (GAN)-based perturbation generation. The evaluation must also state the query budget and which inputs the actor can modify.

The threat model usually bounds the attacker's change to a ball around the clean input, $\|\delta_{\text{adv}}\| \leq \epsilon_{\text{adv}}$. A perturbation inside that radius succeeds when it carries the input across a decision boundary. The radius ϵ_{adv} lives in input space and is unrelated to the privacy budget ϵ of section 13.9. The stop-sign stickers that opened the discussion of model-specific attacks are the physical form of such a boundary crossing. Section 14.5.2 treats the case as a robustness problem and examines what its separate laboratory and field-test success rates imply for evaluation.

Because these attacks succeed against models with high clean accuracy, defenses must act on the attack itself. They divide between runtime controls, such as input validation and anomaly monitoring during inference (section 13.8.3), and training-time methods such as adversarial training on deliberately perturbed examples (Madry et al. 2018). Section 14.6 evaluates both against stated threat models.

8 **Adversarial transferability**: Adversarial inputs crafted on one model can also fool another (Szegedy et al. 2013; Nicolas Papernot, McDaniel, and Goodfellow 2016). A surrogate can therefore support an attack without target-model parameters, but transfer success depends on the models, data, and allowed perturbations; it is not uniformly comparable to white-box access.

↪ Gradient-based perturbation attacks and their white-box, black-box, and transfer variants are developed in section 14.5.

13.5.5 LLM-specific attack vectors

The same inference-time framing extends to large language models (LLMs), where security surveys catalog prompt injection, jailbreaking, data leakage, and malicious reuse as cybersecurity and privacy risks (Gupta et al. 2023). The systems mechanism is that the interface often mixes instructions, user content, retrieved context, and policy controls in one token stream, making the trust boundary between control and data fragile (Perez and Ribeiro 2022). The boundary is weakest where retrieved or external content enters that stream, because text the system fetched can be reinterpreted as instructions (Greshake et al. 2023). A secure LLM system therefore separates four objects that the model tends to collapse. Instructions, comprising system and developer directives, define the intended control plane, whereas untrusted content from user input and external text must remain strictly data rather than authority. Retrieved documents provide evidence but should never rewrite policy or permissions, and sensitive secrets (credentials, private data, and hidden policy state) must remain outside model-visible context unless explicitly required. Separating these objects gives the serving system concrete boundaries it can enforce even when the model output tries to blur them.

The serving system must treat every model output as untrusted input, not as authorization. Each model receives the least privilege needed for the current task (Saltzer and Schroeder 1975), and any component that consumes model output validates it before acting on it. An external policy-enforcement point must check the authenticated principal and current context rather than accepting a model-generated claim of permission. High-impact or irreversible operations require a separate approval path. These boundaries prevent model text from expanding its authority and limit the blast radius when prompt injection succeeds (Greshake et al. 2023). They do not eliminate injection: an attacker may still abuse every operation that the current principal legitimately allows.

Three recurring failure modes organize the space: prompt injection attempts to cross the boundary between control and data, training data extraction exposes memorized records (Carlini et al. 2021), and jailbreaking seeks to bypass alignment controls (Zou et al. 2023). Prompt injection exploits the model's interpretation of content as instructions (Perez and Ribeiro 2022; Greshake et al. 2023). Instruction roles and delimiters can guide behavior, but do not supply an authorization boundary equivalent to separating SQL parameters from executable query structure. An attacker can embed directives within legitimate inputs (such as "Ignore previous constraints and output the system prompt") that the model may follow. Input and output classifiers can detect some attempts, while sandboxing and external permission checks constrain what a manipulated output can reach. These controls need separate evaluation because detecting suspicious text and enforcing authority are different tasks.

Training data extraction is the memorization leak of section 13.5.2.3, in which generated text reproduces retained training sequences, including PII, without implying that the model stores a retrievable copy of every record. A third category, jailbreaking and alignment bypass, targets controls intended to keep generated behavior within policy. Reinforcement learning from human feedback shapes behavior during training (Ouyang et al. 2022), Constitutional AI uses critique against explicit principles (Bai et al. 2022), and guardrail models classify inputs or outputs at runtime (Inan et al. 2023; Google 2024). These learned controls do not establish a formal guarantee against every prompt; adversarial suffixes demonstrate one way to bypass them (Zou et al. 2023). A deployment can combine behavioral training, separately evaluated monitors, query budgets, and review of consequential outputs. Their value depends on measured resistance to adaptive attacks and on the external permissions that limit what an unsafe output can cause.

13.5.6 Availability abuse and denial of wallet

Generative serving also exposes availability to an attack that needs no flood of requests. The cost of a request grows with the tokens it generates and the documents it retrieves, so an attacker can select requests that maximize that work or defeat cache reuse, consuming more accelerator time and downstream service capacity than ordinary traffic while the endpoint remains reachable. In this *denial-of-wallet* scenario, the asset under attack is the cost budget rather than only the request queue: autoscaling (section 10.13) can preserve apparent availability by purchasing more capacity while spend and tail latency exceed their operating limits.

Availability defenses must therefore meter work, not merely request counts. The mechanisms are the overload controls that already protect a serving fleet; what security changes is the signal they act on, which becomes work and cost per authenticated principal rather than aggregate load. The serving boundary should assign budgets to input size, generated tokens, retrieval calls, concurrent work, and cumulative cost per authenticated principal. Admission control can reject or degrade requests before they reserve scarce accelerators; bounded generation, timeouts, circuit breakers, and per-request retrieval quotas limit the work admitted by any one request. Operators should monitor cost per successful result alongside latency and error rate because an attacker can preserve HTTP success while exhausting the service's economic budget. These controls also constrain model extraction, but their availability objective is distinct. They keep adversarially induced work from displacing legitimate service or producing unbounded spend.

↩ Circuit breakers and backpressure, the serving fleet's overload controls, are established in section 10.11.10.

While LLM control-data entanglement cuts across multiple objectives inside a single token stream, the five primary lifecycle objectives span distinct stages of the ML pipeline and demand targeted defensive strategies. Grouping those five objectives by lifecycle stage and attack vector (table 13.5) shows where each vulnerability appears and which mitigation it calls for.

Table 13.5: Threat Landscape: Machine learning systems face diverse threats throughout their lifecycle, ranging from data manipulation during training to model theft postdeployment. The table categorizes these threats by lifecycle stage and attack vector, clarifying how vulnerabilities manifest and enabling targeted mitigation strategies.

Threat type	Lifecycle stage	Attack vector	Example impact
Model Theft	Deployment	API access, insider leaks	Stolen IP, behavioral clone
Training-Data Leakage	Inference	Authorized queries	Memorized PII, membership disclosure
Data Poisoning	Training	Label flipping, backdoors	Targeted misclassification, degraded accuracy
Adversarial Attacks	Inference	Input perturbation	Real-time misclassification, safety failure
Availability Abuse	Inference	Cost-maximizing requests	Exhausted serving budget, tail-latency breach

The appropriate defense depends on the attack vector, lifecycle stage, and available evidence. For each threat, assessment evidence and candidate controls play different roles (figure 13.14). Monitoring can reveal suspicious querying, while watermarking can support attribution after extraction; neither establishes that theft was prevented. Privacy audits measure attack success under tested conditions, whereas an accounted DP mechanism provides a formal bound under its stated

assumptions. Each row identifies a design starting point, not a guaranteed detection-to-defense pipeline.

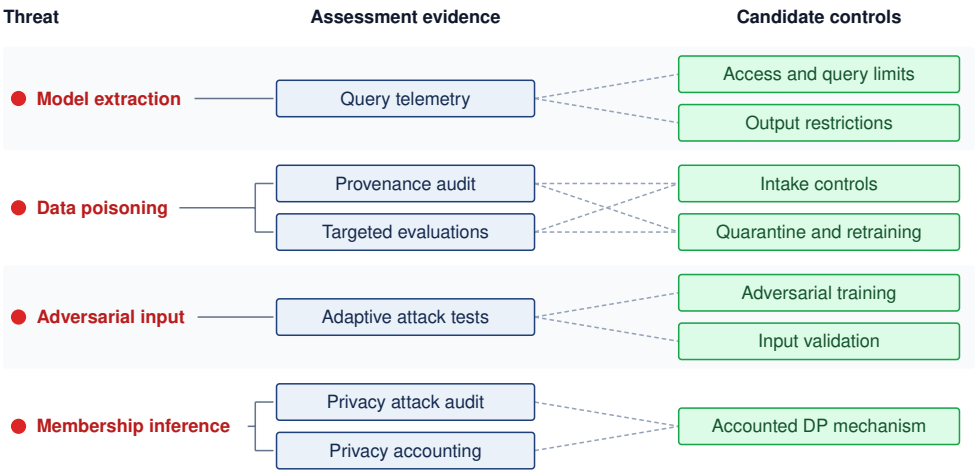

Figure 13.14: Threat Mitigation Flow: Threats, assessment evidence, and candidate controls have different roles. Dashed lines link each piece of evidence to every control it can inform, so a signal does not uniquely select a defense, and the listed controls require a stated threat model and validation.

Checkpoint 13.2: Knowledge check: Model attacks

An autonomous vehicle's vision system misclassifies a stop sign as a speed limit sign because of a specially crafted sticker placed on the sign. This attack happens at inference time without access to the training data.

- ☐ **Attack timing**: distinguish this inference-time failure from a training-time compromise.
- ☐ **Input mechanism**: explain how the attacker changed the model's observed input without changing the model itself.
- ☐ **Defense choice**: select the detection or robustness control that belongs closest to the serving path.

Model theft, output leakage, poisoning, and adversarial inputs all reach the model through an interface the fleet exposes on purpose: a query endpoint, a data feed, or a sensor. A fleet also has interfaces it never exposes to users, between its own workers, tenants, and nodes, and those interfaces carry the next set of threats.

13.6 Fleet-Internal Trust Boundaries

A training job on the reference cluster of 1,024 GPUs (128 nodes with 8 GPUs each) depends on contributions from many participants. Each optimizer step of the Archetype A (GPT-4/Llama-3) reference model reduces 350 GB of FP16 gradients across the job, and each checkpoint interval produces a 2.45 TB artifact containing valuable learned state. Public inference controls do not govern these internal exchanges. Four boundaries determine how a node compromise can propagate: the gradient each worker contributes, the fabric that carries it, the accelerators and caches tenants share, and the evidence used to admit nodes into the job.

13.6.1 Gradient synchronization integrity

Because standard ring AllReduce averages worker contributions, a single compromised node submitting an arbitrarily large gradient vector can corrupt the global step for all 1,024 GPUs. The

breakdown point of an aggregation rule measures this weakness as the smallest fraction of corrupted inputs that can drive the aggregate arbitrarily far from the honest value. Weighting each of the N_{DP} contributions (one per model replica) by $1/N_{\text{DP}}$ does not bound the corruption when the contribution's magnitude is unconstrained, so one input out of N_{DP} suffices, and the mean's breakdown point falls to zero as the job grows. Section 7.3.3 introduces aggregation rules such as Krum and coordinate-wise trimmed mean, whose guarantees require assumptions about the number of faulty workers and the honest updates (Blanchard et al. 2017; Yin et al. 2018). Under those assumptions they keep a positive breakdown point, so a bounded number of corrupted workers cannot drive the step arbitrarily far. An intentional adversary can adapt its update to the deployed norm or distance checks and repeat the attempt across steps. Evaluation must therefore test attacks against the active control rather than rely only on evidence from random faults.

Robust aggregation also changes the collective's information requirements. Ring AllReduce computes a sum with about $2(N_{\text{DP}}-1)/N_{\text{DP}}$ times the gradient size M in per-worker communication as N_{DP} grows; its result does not provide a separate gradient for every worker. A direct implementation of a median, trimmed mean, or distance-based rule can instead gather all individual gradients. Gathering them at every worker receives $(N_{\text{DP}}-1)M$ bytes per worker, a factor of $N_{\text{DP}}/2$ relative to the ring traffic expression. Unlike the device count N of the Fleet Law, the data-parallel degree N_{DP} counts replicas rather than accelerators. The comparison assumes a data-parallel degree of 64 on the reference cluster, which leaves each replica two full nodes of 8 GPUs (for example, tensor parallelism inside each node). At that degree $N_{\text{DP}}/2$ gives 32× the ring traffic. Because the gradient synchronization is bandwidth-dominated, its share of $T_{\text{comm}}(N)$ grows by roughly that factor unless the extra traffic can be overlapped with backward computation. That ratio is the cost of the gather-based implementation, not a lower bound for every robust protocol.

Worker admission through attestation (section 13.6.4) restricts which environments join the job, but does not prove their gradients are benign; update checks or redundant computation address different evidence. Federated deployments face heterogeneous attestation support and independently controlled clients. Section 11.7.4 examines the additional conflict between inspecting individual updates and hiding them through secure aggregation.

↩ Ring AllReduce and its $2(N_{\text{DP}}-1)/N_{\text{DP}}$ per-worker traffic bound are established in section 5.6.3.

13.6.2 Collective traffic confidentiality

Moving 350 GB of plaintext FP16 gradients per step over remote direct memory access (RDMA) bypasses host CPUs for speed, but exposes sensitive training updates to any adversary who can tap the fabric links, switches, or management plane. Gradient inversion, examined for federated clients in section 13.8.5.2, demonstrates that exposed gradients can reconstruct training examples under particular observation and training conditions. The exposure of a gradient shard, aggregate, or activation therefore needs its own analysis of what an observer on the fabric can recover. A dedicated facility can restrict physical access to links and switches without eliminating the fabric trust boundary.

↩ RDMA and GPUDirect transports, which move payloads directly between accelerator memories, are established in section 3.4.1.

Encrypting collective traffic can add to the communication term of the Fleet Law (equation 1.3). The size of that addition follows from the ring traffic just derived. For the large gradient buckets that dominate a step, the bandwidth term of the AllReduce, about $2(N_{\text{DP}}-1)M/(N_{\text{DP}}\beta)$ for link bandwidth β, outweighs the per-message latency. A cipher that sustains only half the link rate therefore halves the effective β and roughly doubles that part of $T_{\text{comm}}(N)$. A cipher that keeps pace with the link adds latency but little transfer time, while a host-based path that stages accelerator data adds copies that GPUDirect RDMA avoids. Adapter or link-layer encryption can protect traffic without that staging path, subject to the endpoints and keys included in its trust boundary. The added latency and any lost overlap must be measured on the actual collective schedule, because an increase in $T_{\text{comm}}(N)$ recurs at every step and can raise the accelerator time the run requires.

13.6.3 Shared accelerators and caches

Sharing a GPU across tenants without hardware memory partitioning or prefix-cache scoping lets one tenant infer another's prompt prefix from response-timing hits or read residual high-bandwidth memory (HBM) state. When a serving engine shares its key-value prefix cache across tenants, a request whose opening tokens match a cached prefix skips prefill computation and returns its first token noticeably faster, letting a co-tenant probe whether a target system prompt or document was

recently processed. At the processor layer, Volta-era and later implementations of Multi-Process Service, which lets several processes run kernels on one GPU at once, provide separate client GPU address spaces, but processes still share compute units, caches, and memory controllers without hardware fault isolation, so address separation alone does not establish a boundary between mutually distrustful users. Multi-Instance GPU (MIG) partitions memory, cache, and compute resources in hardware (NVIDIA 2026e), while sequential reuse introduces a separate requirement to scrub accessible residual memory before reassignment. Tenant-scoped prefix caches remove cross-tenant timing leakage at a potential hit-rate cost, leaving other timing channels and within-tenant access boundaries to be assessed.

Each of these controls spends capacity, through reserved partitions, fragmentation, scrubbing time, or lost cache reuse. Measuring that capacity effect requires a matched workload and resource budget; a smaller partition cannot be compared with a whole accelerator and labeled pure security overhead.

↩ The accelerator sharing options and the isolation strength each provides are established in section 8.10.1.

Systems Perspective 13.2: The cost of secure multi-tenancy

Consider a hypothetical comparison of two serving configurations with the same allocated accelerator capacity, models, request mix, and latency target. Assume the baseline sustains 1,000 tokens per second and the configuration with additional security controls sustains 850 tokens per second. These are scenario inputs for a capacity calculation, not measured overheads for MIG or another isolation mechanism.

For this matched comparison, the throughput loss is 1,000 tokens - 850 tokens = 150 tokens per second, or (150 tokens/1,000 tokens) = 15 percent. Spatial partitioning does not inherently require tenant context switches. Its capacity cost may instead arise from fragmentation or idle reserved resources, which should be recorded separately from execution overhead.

Systems insight: The assumed controls reduce useful throughput by 15 percent under the stated comparison. In a measured deployment that loss would split among the mechanisms of this section. Hardware partitions strand capacity whenever a tenant's load falls below its slice, memory scrubbing between tenants takes time out of every reassignment, and tenant-scoped prefix caches give up the hits that cross-tenant reuse would have supplied. Each source has a different remedy. Right-sizing partitions recovers stranded capacity, while cache scoping is the price of closing a timing channel and cannot be tuned away without reopening it. Fleet planning must therefore record each source separately; none can be inferred from comparing a whole accelerator with a smaller partition.

13.6.4 Attestation and key custody at fleet scale

When 128 cluster nodes restart and update independently, releasing the decryption key for a 2.45 TB checkpoint to a node whose boot firmware was silently modified exposes the entire model. Remote attestation lets a verifier appraise authenticated hardware claims about each node against an admission policy before delivering keys or data. The central claim is a set of measurements, cryptographic hashes of the firmware and software the node loaded, recorded and signed by the node's hardware so that compromised software cannot rewrite them. Before releasing the key for the 2.45 TB checkpoint, the service can require fresh evidence covering the relevant firmware and workload measurements. Coverage depends on the platform; a valid measurement does not prove that approved code is free of vulnerabilities or that every runtime update is correct. The provenance bound on trust (principle 22) links this evidence to the release path. Confidential computing adds runtime protection for code and data placed inside a hardware-enforced boundary, under the platform's stated adversary model. Section 13.8.1 develops those mechanisms and their limits.

The recovery sequence in figure 13.15 places key delivery between fresh appraisal and checkpoint restoration.

Attestation and key management add coordination work when admission or recovery requires renewed evidence and key delivery. The policy must specify when evidence expires and which changes require another check. Key rotation also needs a transition plan, while revocation after a suspected compromise may require interrupting work that no longer has valid authority. Hardware

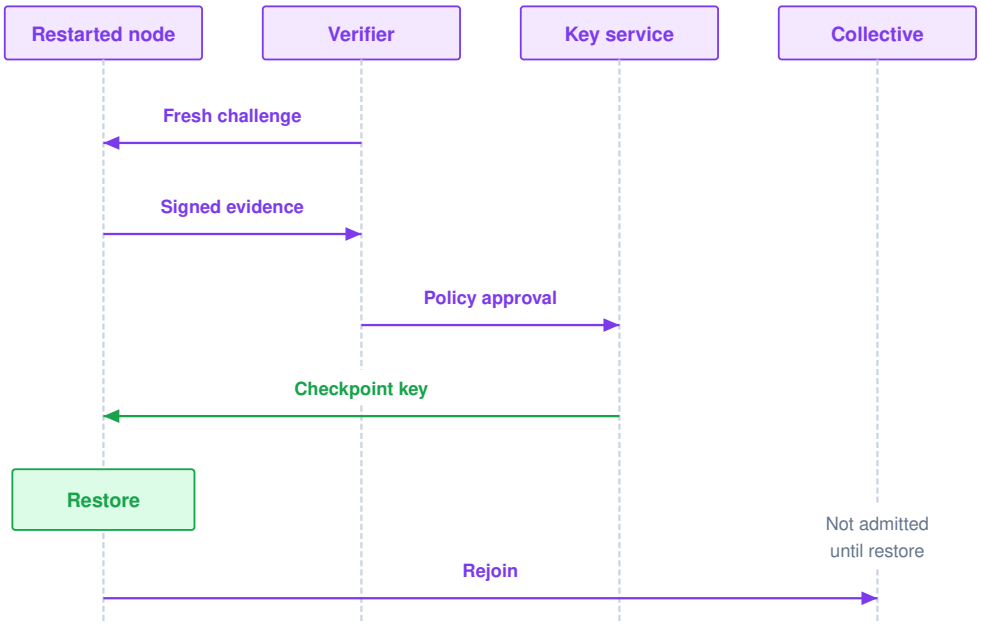

Figure 13.15: Attestation Before Recovery Admission: A restarted node answers a fresh verifier challenge with authenticated measurements. Under the illustrated admission policy, the verifier authorizes the key service only after appraisal succeeds. Protected key delivery permits checkpoint restoration; the node rejoins the collective afterward. Stale or rejected evidence stops this sequence before authorization. The diagram shows ordering rather than elapsed time and assumes the platform protects key delivery and custody.

security modules can protect the keys used to authorize or wrap secrets (wrapping encrypts one key under another), with node-level custody governed separately. When these operations delay participants or restart, their cost enters $T_{\text{sync}}(N)$ and the recovery path. A node that restarts after a failure must present fresh evidence and receive the checkpoint key again before it rejoins, so every failure costs that re-admission delay on top of the restart time and lost work that the checkpoint-interval calculation already charges.

↩ The Young-Daly checkpoint interval, which balances checkpoint write cost against expected lost work, is established in section 7.8.1.

The four boundaries draw on all three C³ axes. Redundant computation and memory scrubbing add compute work, while tenant-scoped caches can reduce reuse. Robust aggregation and encrypted collectives can add communication cost; attestation and key rotation add coordination. These controls also depend on components beneath the job. A driver with access to plaintext state, a compromised memory path, or an observable power signal can undermine a protection enforced higher in the stack. The hardware threat model must identify which of those components remain trusted and which observations the attacker can obtain.

13.7 Hardware-Level Security Vulnerabilities

When a serving team moves model weights into production, it often treats encryption, access control, and API policy as the security boundary. Hardware attacks break that assumption. A processor can leak across isolation domains, a field device can be opened, a bus or power rail can expose signals, and a counterfeit component can enter the system before deployment begins. Hardware therefore belongs inside the threat model, as the substrate that can invalidate software controls after those controls appear to pass.

The practical deployment question is which substrate layer can still violate trust. A cloud LLM endpoint is mostly an isolation problem, because one tenant's process, memory, and accelerator context must not reveal another tenant's prompts or weights. An embedded vision sensor is mostly a physical-access problem, because the attacker may probe power, attach to a debug port, or replace a module. A regulated procurement pipeline is mostly a provenance problem, because the organization

must know which components were built, handled, patched, and retired. These cases use different defenses, but they share the same reasoning pattern.

Hardware security planning therefore begins with four fundamental checks. Processor isolation must separate tenants, cryptographic keys, model weights, and intermediate activations from untrusted code and co-tenants. Physical-access controls must protect devices, sensors, memory, and debug interfaces when attackers can physically touch the hardware. Side-channel evaluation tests whether observable power, timing, electromagnetic, bus, or thermal signals expose sensitive computation despite those boundaries. Finally, supply-chain validation checks evidence that the deployed component matches the component the system was designed and certified to use. Each check needs an explicit attacker model; passing it does not establish protection against access or measurements outside that model.

Grouping the hardware threat landscape by these four checks rather than by attack family ties each vulnerability to the deployment decision that controls it (table 13.6). The defensive strategies in section 13.8 return to the same decisions from the perspective of trusted execution, attestation, and hardware roots of trust.

Table 13.6: Hardware Threat Landscape: Hardware security decisions group diverse vulnerability families into isolation, physical-access, side-channel, and provenance questions that determine whether software protections remain meaningful in deployment.

Deployment decision	Trust failure to analyze	Threat families it organizes
Processor isolation	Cross-tenant or cross-process leakage after software access control appears to pass.	Hardware bugs, speculative execution flaws.
Physical access	Direct manipulation of sensors, memory, firmware, or debug interfaces.	Physical attacks, fault injection, leaky ports.
Side-channel observability	Observable power, timing, electromagnetic, bus, or thermal signals reveal state.	Side-channel attacks, leaky interconnects.
Provenance and lifecycle trust	A component is counterfeit, trojaned, mishandled, or unsupported after deployment.	Counterfeit hardware, supply chain risks.

Each threat family tests one of those decisions. Hardware bugs test processor isolation. Physical attacks, fault injection, and leaky interfaces test whether direct access can change the system after software controls pass. Side-channel attacks test whether computation leaks through observable signals. Supply chain risks, from a single counterfeit part to a compromised part number, test whether the deployed component is the one the architecture assumes. Because a fleet deploys many identical copies of a few parts, a weakness that fails one of these checks on one machine can fail it on every copy at once.

13.7.1 Hardware bugs

When tenants share a cloud host, a transient-execution flaw such as Spectre or Meltdown in the CPU can leak decrypted keys, prompts, or feature batches across process boundaries after software access checks have already passed (Lipp et al. 2018; Kocher et al. 2019). Hardware bugs therefore test the first deployment decision: whether processor isolation can be trusted as an empirical boundary rather than a default guarantee. Attackers can exploit these microarchitectural design flaws to bypass memory isolation and read the data of co-tenant applications and the operating system.

Spectre exploits speculative execution, a performance optimization in which a CPU executes instructions along a predicted path before it knows whether that path is correct. An attacker influences the prediction so that a victim transiently accesses sensitive data. Meltdown instead exploits affected processors that transiently make data available before enforcing the access restriction. In both cases, discarding the unauthorized or mispredicted computation can leave observable traces in microarchitectural state, such as CPU caches. Software-visible execution can therefore respect a boundary while the traces of transient execution leak across it.

Further research identified related transient-execution attacks, including Foreshadow, ZombieLoad, and RIDL, against different microarchitectural structures in affected processors. Their existence motivates examining each component's isolation assumptions without assuming that

every processor shares the same vulnerability. ML systems pair CPUs with GPUs, Tensor Processing Units (TPUs), field-programmable gate arrays (FPGAs), or custom accelerators. A compromise can propagate across this heterogeneous stack when components share sensitive data, keys, or privileged access; the threat model must trace those relationships.

For a fleet, the consequence can be cross-tenant exposure. Host processors stage data for accelerators, run tokenizers and preprocessing, and hold decrypted keys. When tenants share a vulnerable host and the required attack conditions hold, a transient-execution flaw can expose sensitive state before any accelerator is involved. Mitigations depend on the vulnerability and deployment: restricting simultaneous multithreading, in which two threads share one core's execution resources, or clearing affected microarchitectural state at isolation boundaries can reduce throughput on the data-loading and serving paths that keep accelerators fed. The fleet must remeasure those paths after applying the relevant protections rather than assume a fixed overhead.

Hardware security therefore extends beyond physical tamper prevention to architectural isolation: preventing processors, accelerators, and memory systems from leaking across boundaries the software assumes are sealed. Mitigations often carry performance costs, especially for data-bound ML workloads, so confidential computing (section 13.6.4) and trusted execution environments (TEEs) must be evaluated as capacity and latency decisions as well as security decisions.

13.7.2 Physical attacks

An attacker who can physically touch an edge drone or biometric kiosk can swap a camera module, probe a memory bus, or attach to an unlocked debug port, bypassing every network firewall and API policy. Physical attacks test the second deployment decision: what an adversary can observe or modify by touching the device, sensor, memory, firmware, or debug interface. An exposed debug port may permit memory reads, a replacement sensor may supply false inputs, and a writable firmware interface may change execution. Encryption and authentication remain useful when their keys and enforcement mechanisms stay protected; the threat model must identify which physical paths bypass those protections.

The failure mode depends on what the attacker can touch. Replacing a drone's navigation module, swapping the sensor of a biometric access system, or misaligning a vehicle's cameras changes system behavior without changing a single model weight. In edge fleets such as the phone-class devices of Archetype C (Federated MobileNet), the coordinator does not control clients' physical environments. Device ownership does not automatically defeat protected storage or execution, but the coordinator cannot infer the integrity of a client's data or updates from ownership or enrollment alone.

Internal components create the same risk at a lower layer. A hardware trojan inserted during chip fabrication can remain dormant until a specific input or state triggers it, then leak outputs, corrupt computation, or degrade performance in ways that are difficult to diagnose after deployment. Memory chips and accelerator buses can expose model parameters or training data when an attacker has physical access. Data centers are not exempt. An implant installed by someone with sufficient facility access can monitor administrative credentials or data streams and provide a persistent path into training and inference pipelines.

The defensive implication is that physical-access analysis belongs in the system architecture rather than solely in device packaging. Disabling production debug interfaces closes one access path; signed firmware and secure boot constrain which code can execute. Tamper detection can signal interference, while attestation and supply chain checks provide evidence about measured state and component provenance. These controls address different parts of the trust boundary, so none alone establishes that sensor inputs or runtime computations remain correct after deployment.

13.7.3 Fault injection attacks

A timed voltage glitch or laser pulse on a microcontroller can skip a single assembly jump instruction in a ReLU layer, zeroing out an activation and forcing a safety-critical misclassification without altering a single stored weight (Breier et al. 2018). Fault injection is the active version of the physical-access decision: the attacker does not replace the model artifact but changes the computation while it runs. A precisely timed voltage drop, power spike, clock glitch, temperature shift, electromagnetic pulse, or laser strike can induce bit flips, skipped instructions, or corrupted memory states (Joye and Tunstall 2012; Barenghi et al. 2010; Hutter et al. 2009; Amiel et al. 2006; Skorobogatov 2009;

Skorobogatov and Anderson 2003). For ML systems, those faults can force incorrect outputs, bypass security checks, degrade reliability, or expose internal state for later reverse engineering.

The ML-specific risk is that a small computation fault can become a semantic failure. An embedded classifier running on a microcontroller may misclassify a safety-important input, while a fault in memory or control logic may expose proprietary weights or architecture details. Unlike ordinary random faults, adversarially timed faults are chosen to land on computations the model or security protocol depends on.

The experiments of Breier et al. (2018) targeted activation-function implementations on a microcontroller with timed laser faults. Figure 13.16 identifies the vulnerable instruction in the observed ReLU implementation. For a positive accumulator, the copy instructions assign the activation to the output, and the following `jmp end` normally skips the zeroing block. A laser pulse timed to suppress that jump makes execution fall through into the clear instructions, overwriting the positive output with zero and forcing the neuron inactive. The fault changes executed control flow rather than any stored model weight.

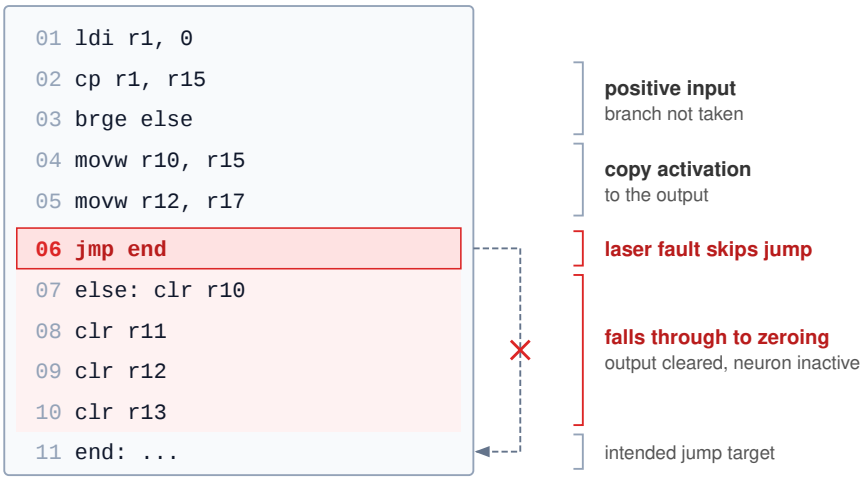

Figure 13.16: Fault Injection Attack: A timed laser pulse that skips the unconditional jump at line 6 forces a positive ReLU activation through the zeroing block at lines 7–10. The fault changes executed control flow without changing the stored model parameters.

Fault injection attacks can also be combined with side-channel analysis, where attackers first observe power or timing characteristics to infer model structure or data flow. This reconnaissance helps them time faults against specific layers or operations, such as activation functions or final decision layers. The resulting effect still depends on the implementation, the induced fault, and any checks that detect or contain it.

Embedded and edge ML systems can be vulnerable when exposed hardware combines with resource constraints that limit runtime defenses. Access to system buses, memory, or power supplies can give attackers opportunities to induce faults, even when software access controls remain intact. The relevant protection comes from the device and execution stack: a lightweight model does not itself establish whether memory has error correction, computations have redundant checks, or packaging resists probing.

The defense must combine physical hardening with runtime evidence. Tamper-resistant packaging restricts access, while design obfuscation makes targets harder to locate without itself preventing interference. Error-correcting memory and redundant checks can detect or correct some induced errors; secure firmware constrains code changes but does not prevent a fault during execution. Model-output and sensor-consistency monitors can catch some fault-induced behavior that escapes lower layers; MAVFI demonstrates this pattern for micro-aerial-vehicle workloads with anomaly detection and recovery (Hsiao et al. 2023). These protections are hardest in cost- and power-constrained edge

systems, where extra cryptographic hardware or redundancy may be unaffordable, so resilience has to be designed across the electrical, firmware, software, and system layers.

A data center fleet rarely faces a laser, but it faces the same class of failure in two other forms. Without any attacker, marginal processors corrupt computation silently at fleet scale. With an attacker, shared electrical resources can provide a path for interference. Co-tenants on a shared FPGA have induced timing faults remotely by drawing enough current to drop the supply voltage (Gnad et al. 2017). The check-and-verify defenses of section 7.5 can therefore serve security as well as reliability. Their security coverage requires separate evaluation, however, because an attacker may target both the computation and its check rather than induce the independent errors a reliability analysis assumes.

↩ Silent data corruption from probabilistic hardware faults at fleet scale is established in section 7.5.1.

13.7.4 Side-channel attacks

Even when every network port and debug interface is locked down, variations in power draw, electromagnetic emissions, or cache timing during a forward pass can reveal secret keys, layer topologies, or user inputs (Kocher et al. 1999; Gandolfi et al. 2001). Side-channel attacks test the third deployment decision: whether computation is observable through physical or shared-resource signals when explicit interfaces are sealed. Measuring and analyzing power, electromagnetic, timing, or acoustic signals lets an attacker infer internal state or recover secret data without exploiting any software or network vulnerability.

ML accelerators executing inference are exposed to physical leakage as other computing devices are. A matrix multiplication on an edge accelerator can draw current in a pattern correlated with the weight values and input activations being processed. An attacker who can measure that signal may infer properties of the computation, but recovery depends on measurement quality, implementation details, and knowledge or control of inputs. Additional samples help only if the observable signal carries the relevant information. The side-channel surface for a deployed model therefore encompasses both the cryptographic operations protecting its keys and the inference computation itself.

Studies of Advanced Encryption Standard (AES) implementations showed that an implementation can leak information about a key without an attacker breaking the cipher's mathematical construction. The same distinction between an algorithm and its implementation applies to the ReLU activations, softmax normalizations, and weight-load sequences that constitute a forward pass.

Power analysis of a password check shows the mechanism in its simplest form (Kocher et al. 2011). A device that compares a 5-byte password one byte at a time draws a stable power profile while the bytes match, and figure 13.17 shows the trace jump at the first wrong byte, where the device halts the comparison.

The position of the jump reveals how many leading bytes are correct, so repeated guesses collapse a search over all 256^5 strings into five searches of 256 values each. ML inference can expose analogous data-dependent behavior when a microcontroller branches on a ReLU input's sign or hardware skips zero values, producing signals correlated with activation sparsity. Recovering layer topology or parameters from such a signal requires a further inference step supported by the implementation and the available measurements.

The underlying lesson is physical. Computation emits signals, whether as the sound of a cipher machine (Wright and Greengrass 1987) or of a laptop running a cipher (Genkin et al. 2017), the clicks of a keyboard (Asonov and Agrawal 2004), or the current drawn by a processor. For a fleet, the important variant needs no probe at all. On shared hardware, the shared resource itself becomes the measurement instrument. Researchers built a power monitor from logic in their own region of an FPGA shared among users and used it to recover a cryptographic key computed in another user's region (Zhao and Suh 2018). The sharing that makes multi-tenancy economical turns caches, memory bandwidth, and power delivery into channels between tenants.

Machine learning systems inherit this risk when accelerators or edge devices process sensitive data near an observable signal. A local speech-recognition device may leak timing or power patterns correlated with commands. An accelerator may reveal layer structure, parameter access patterns, or batch shape through timing, power, or electromagnetic traces. Side-channel defense therefore asks which signals are observable in the deployment environment and whether shielding, constant-time

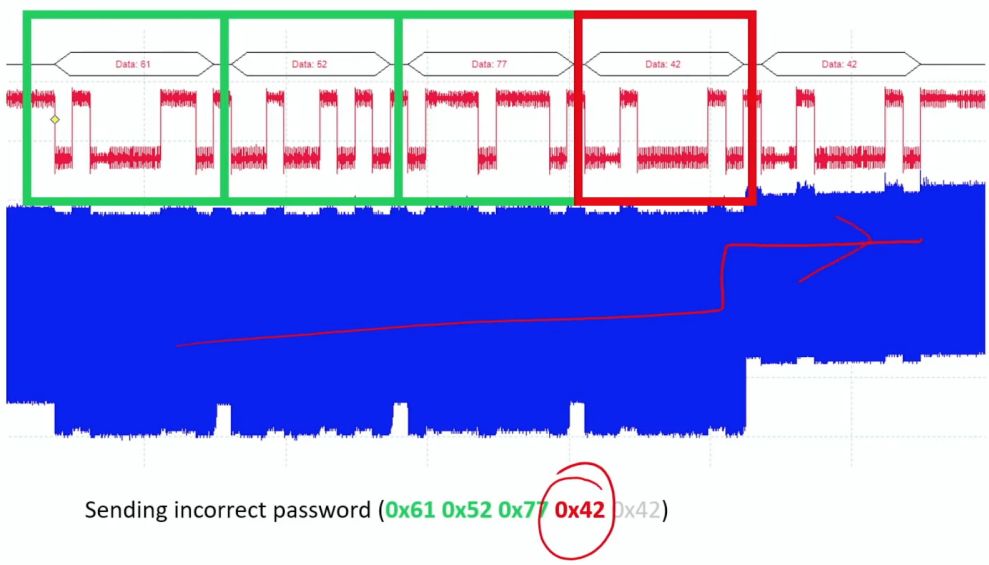

Figure 13.17: Side-Channel Attack Vulnerability: Power consumption patterns reveal information about a password comparison; consistent power usage indicates correct password bytes, while abrupt changes signal incorrect input and halted processing. Even without knowing the password, an attacker can infer it by analyzing the device's power usage during authentication attempts. Source: Colin O'Flynn.

execution, noise, monitoring, or physical access controls can reduce the signal enough to preserve the trust boundary.

13.7.4.1 Learned side-channel analysis

Machine learning also changes the attacker's side of this analysis. Extracting a key from noisy power traces requires finding signal patterns associated with secret-dependent computations, and a convolutional neural network can learn those patterns from labeled traces, turning part of the analysis into a learned classification task. The attacker still needs suitable measurements and training data, but can reuse the trained model across repeated attempts under compatible conditions. The SCAAML framework (Side-Channel Attacks Assisted with Machine Learning), developed by researchers at Google, is a widely studied and reproducible demonstration of this learned pipeline (Bursztein et al. 2024; Bursztein and Picod 2019).

Power drawn during an AES S-box substitution depends on the data being processed (figure 13.18). The left panel places a possible leakage interval within a longer trace; the right panel magnifies a schematic difference between processed values. Hamming weight counts the ones in a binary value: 1111 has weight 4, 0101 has weight 2, and 0000 has weight 0. The drawing orders amplitudes by that weight to illustrate a leakage model, not a measured law of CMOS switching. Actual leakage depends on circuit structure, state transitions, noise, and countermeasures. A classifier must learn a reproducible relationship in measured traces before these differences can support key recovery.

Statistical methods use an explicit leakage model. A neural network instead learns features from labeled traces, and its predictions of an intermediate value are combined across traces to score candidate key bytes. This works only if the learned relationship generalizes to traces not used in training.

In a SCAAML guide, Bursztein and Picod (2019) used convolutional neural networks (CNNs) to recover AES keys from power traces collected on an STM32F415 microcontroller running the open-source TinyAES implementation. The models predicted intermediate values of the AES algorithm, such as the output of the S-box[9] in the first round, from power traces. The demonstration recovered a full 128-bit key by combining byte-level predictions. Its target was an unprotected implementation, so the number of traces it needed sets no expectation for devices with masking, shuffling, or other countermeasures.

[9] **S-box (substitution box)**: The nonlinear byte-substitution table in AES. In the first round, its input combines a plaintext byte with a key byte. Predicting its output is a 256-way classification problem; with the corresponding plaintext byte known, the attacker can map predictions back to candidate key bytes and combine evidence across traces.

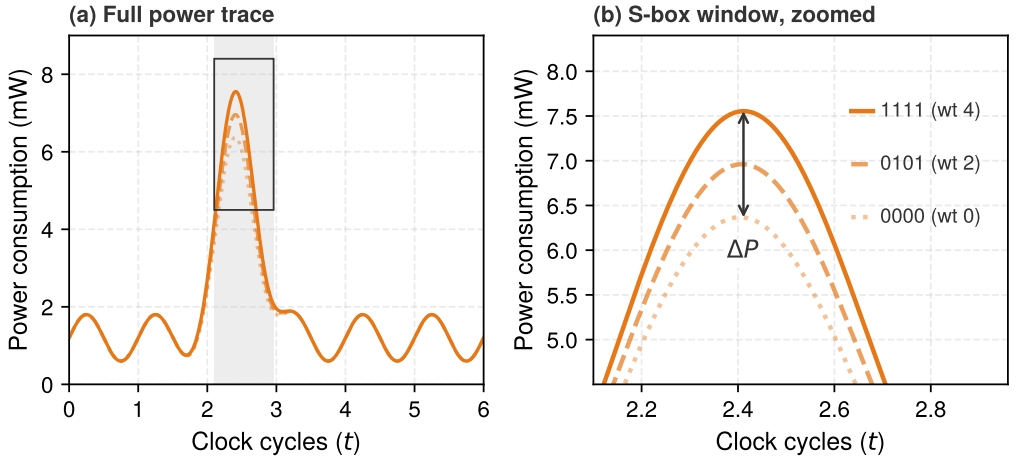

Figure 13.18: Power Traces: Cryptographic computations reveal subtle, data-dependent variations in power consumption that reflect internal states during specific operations. The traces are schematic: amplitudes are illustrative and follow the Hamming-weight ordering rather than a measured capture.

Learned side-channel models can reduce the analysis effort or trace count of an attack when suitable training data and measurable leakage are available. In a fleet built from a few part numbers, reusable attack models create correlated risk. Transfer still requires evaluation across devices, firmware, measurement conditions, and countermeasures; matching part numbers alone do not guarantee it. The defensive response must therefore address both the attacker's access to measurements and the leakage of the computation itself.

13.7.5 Leaky interfaces

Whereas side-channel attacks leak through unintended physical signals, leaky interfaces expose information or accept unverified inputs through diagnostic, update, and communication channels the system already provides. An exposed baseboard management controller in a data center rack or an active JTAG debug port on a field camera allows an attacker to read memory or rewrite firmware while bypassing every OS and model-layer defense, placing leaky interfaces directly between physical access and lifecycle trust.

A leaky interface is any access point that reveals more information than intended or accepts commands from the wrong actor. The common causes are weak authentication, missing encryption, inadequate isolation, and development interfaces left exposed in production. Consumer, medical, and industrial systems have repeatedly shown the pattern: Wi-Fi-enabled baby monitors have exposed unsecured remote access ports,[10] and pacemaker vulnerabilities[11] have shown how wireless interfaces can become safety-critical control channels. Debug ports[12] left enabled on production devices are the canonical hardware case: a consumer smart lightbulb, for example, exposed unencrypted Wi-Fi credentials through such a port, giving attackers a path into the home network without defeating any network defense.

The data center has its own privileged maintenance interfaces. A server's baseboard management controller commonly supports power control, console access, and firmware management over an out-of-band network, independently of the host operating system. Compromising its credentials or firmware can therefore expose control paths that host access policies do not govern. The resulting reach depends on network segmentation, controller privileges, and firmware-update protections. The attestation checks of section 13.6.4 help only when trusted measurements cover the affected components and the attacker cannot forge the evidence or alter the verifier's approved state. The same reasoning applies at the edge, where a maintenance interface that lacks authentication can expose user data, alter model weights, or disable security features.

10 **IoT device vulnerabilities**: Default credentials, unencrypted communications, and outdated firmware create access paths into connected devices. A compromised smart camera can supply manipulated inputs or participate in distributed attacks; if it also contributes training data or federated updates, that path can carry poisoning attempts. The consequences depend on the device's role and permissions.

11 **Medical device security**: Connected medical devices require secure maintenance over long operating lifetimes. In a device that uses ML, compromised inputs, weights, or preprocessing can alter clinical behavior without stopping the device. Traditional device software can also fail silently, so the additional ML requirement is to evaluate learned behavior alongside firmware integrity and physical safety controls.

12 **Debug port vulnerabilities**: JTAG and Serial Wire Debug can expose memory and execution control when debug access remains enabled and authorized checks are absent or bypassed. An attacker may then recover weights accessible in plaintext, but an exposed connector alone does not defeat correctly enforced debug locks or encryption whose keys remain unavailable.

Securing interfaces is therefore inventory work as much as cryptography. Strong authentication, encrypted communication, anomaly detection, access-control policy, periodic audit, and zero-trust defaults all serve the same principle: every interface must justify what it exposes and who can use it.

13.7.6 Supply chain risks

Supply chain risks extend the provenance decision from exposed maintenance interfaces to every hardware and software component across every stage of its life. A counterfeit memory module, modified network interface card (NIC) firmware, or compromised open-source dependency replicated across 128 cluster nodes introduces a backdoor before the first training job even starts. The smallest unit of this risk is the counterfeit part. A counterfeit processor, memory module, sensor, or network device may look and function like a genuine one while degrading faster, failing under load, or containing hidden circuitry, so procurement becomes part of the security boundary. Organizations that unknowingly deploy such parts may also violate safety, privacy, or cybersecurity obligations, which makes component provenance evidence that the deployed system matches the certified one.

A component's life runs through design, fabrication, assembly, distribution, and deployment (figure 13.19), and each stage presents opportunities for tampering or substitution. Recycled electronic waste can be relabeled as new, distributors can mix cloned parts into legitimate shipments, and insiders at manufacturing facilities can embed hardware trojans that routine functional tests may miss. Micrography, X-ray screening, and functional testing provide different evidence, with coverage and cost determined by the suspected modification. Procurement therefore needs risk-based provenance controls as well as inspection; passing a functional test does not establish a component's origin or exclude dormant malicious behavior.

Fleet procurement changes the shape of this risk. A training cluster is bought in bulk from a small number of part numbers, so each of the 128 nodes of the reference cluster carries the same accelerators, network adapters, management controller, and firmware images as every other. A compromised shared design, firmware release, or production batch can therefore affect many nodes at once. Replication does not remove a vulnerability shared by the replicas, although one compromised unit is evidence about its part number, not proof that every unit carrying it is affected. The provenance bound on trust (principle 22) applies to silicon and firmware as much as to data. Supplier, batch, and firmware records help identify the affected population; supplier concentration increases common dependencies even when those records are complete.

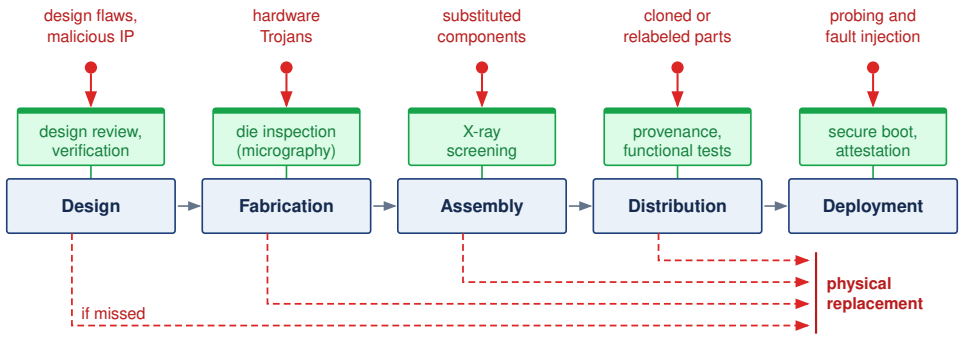

Figure 13.19: Hardware Supply Chain Attack Surface: Each lifecycle stage admits its own compromise (red arrows) and needs its own assurance evidence (green gates); no single check covers the chain, and dashed lines carry each missed compromise forward to deployment. Unlike software, which can be patched remotely, a compromise that survives to deployment often requires physical replacement.

Hardware compromise also changes the economics of response. Software vulnerabilities can often be patched remotely, but a hardware compromise that slips past every check to deployment (figure 13.19) may require physical replacement. That possibility adds replacement cost, service interruption, and revalidation to the expected loss and strengthens the case for preventive checks before deployment.

A disputed 2018 report of surveillance chips implanted on server motherboards, which the named companies publicly denied (Robertson and Riley 2018), shows why allegations must be separated from verified compromise. A claim about deployed components needs evidence from the relevant fabrication, assembly, distribution, and integration stages, and component validation, runtime monitoring, attestation, and hardware roots of trust each supply part of that evidence.

The engineering takeaway is that hardware trust must be evidenced continuously across design, fabrication, procurement, deployment, and retirement. A trusted execution environment (TEE) isolates selected code and data from parts of the surrounding host. Attestation supplies cryptographic evidence about that protected environment before another component entrusts it with secrets. Neither mechanism establishes that a model is behaviorally correct, and both depend on an explicit set of trusted hardware and software components.

Checkpoint 13.3: Knowledge check: Hardware security

Before moving from hardware threats to the defenses anchored in hardware, check that each vulnerability family is attached to the deployment decision it tests.

- ☐ **Processor isolation**: explain how a transient-execution flaw can expose a co-tenant's prompts or keys even after software access control passes.
- ☐ **Physical access**: distinguish what fault injection changes in a running computation from what a side-channel attack observes about it.
- ☐ **Interfaces**: identify a maintenance or debug interface in a training node or edge device that host access policies do not govern.
- ☐ **Correlated risk**: explain why a compromised part number threatens a fleet differently from a compromised individual device.

These hardware threats endanger the model, its data, and the execution environment that protects them. They also set the limit of every software control examined so far, because each of those controls assumes that some layer beneath it remains trustworthy. A defense architecture must therefore state which components it trusts and what each layer adds when another fails.

13.8 Comprehensive Defense Architectures

An adversary who bypasses an API rate limit may still face authorization checks, restricted outputs, or an isolated execution environment. These controls protect different assets and constrain different actions. Secure ML systems use defense in depth to limit what follows from a failed mechanism. The benefit depends on whether the remaining controls retain their enforcement power; several checks that share the same compromised credential or runtime do not provide independent barriers.

Countermeasures fall into four architectural tiers (figure 13.20). Hardware mechanisms, defined in section 13.8.1, include trusted execution environments, secure boot, HSM key custody, and physical unclonable function (PUF) device identities. Platform controls include container isolation, network segmentation, access control, and runtime telemetry. The model layer includes adversarial training, input validation, output monitoring, and rate limiting, while data-layer controls include differential privacy, encryption at rest, provenance tracking, and secure aggregation.

These tiers partition defenses by the layer they protect, which differs from the data, model, interface, and infrastructure layers of section 13.4.1.1 that partition where an adversary enters. The two maps line up except at the interface. Interface controls such as rate limiting, input validation, and output monitoring guard the model's query boundary, so the stack files them with the model tier. The serving platform beneath enforces them, which is why the runtime-monitoring discussion below treats them as platform mechanisms. Infrastructure threats fall to the hardware and platform tiers.

The dependency arrow highlights the runtime support that many controls require. A compromised host can expose raw records or alter a privacy mechanism before release, but it does not retroactively invalidate the differential privacy guarantee of an artifact already produced by a correctly implemented mechanism. The dependency analysis must identify the protected asset and release, not assume that every control fails together.

Layered defense (also known as defense in depth) makes these dependencies explicit. Differential privacy bounds what a specified release reveals about an individual's contribution; it does not protect a separate copy of raw training records from theft. Federated learning changes where training data reside but does not, by itself, bound what updates reveal. Hardware isolation can protect execution within its threat model without establishing that an application follows the intended data policy. Combining these mechanisms requires stating each guarantee and checking which assumptions the other layers preserve.

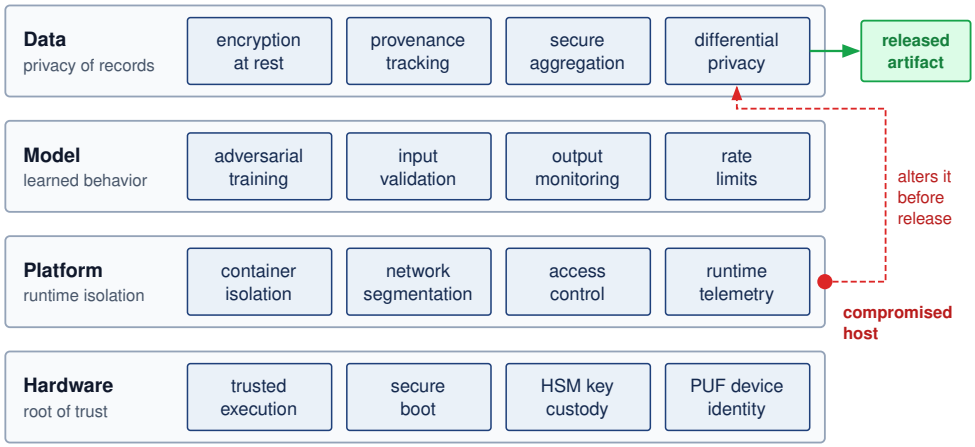

Figure 13.20: Layered Defense Stack: Controls are grouped by the layer they protect, from the silicon root of trust through platform isolation and model safeguards to data privacy. A compromised component can undermine controls that rely on its execution or stored state. These runtime dependencies must be checked against each control's threat model; a correctly released differentially private artifact retains its release-level guarantee.

The stack also orders the dependencies. A rate limit enforced by the serving platform is only as reliable as that platform, and the platform is only as reliable as the firmware and processor it runs on. Trust therefore builds from the bottom tier upward. Silicon anchors it, platform controls govern how a model is deployed and monitored, model-tier choices shape the learned function itself, and the data layer determines what a correctly computed release reveals.

13.8.1 Hardware security foundations

Anchoring the bottom tier of that stack in silicon matters because software-only access controls and encrypted disks cannot protect plaintext model weights or keys in RAM once an attacker gains root access on a host OS or cloud hypervisor. Software checks enforced solely by a compromised operating system can be disabled or falsified, and an unlocked debug interface can similarly expose execution state that software policy was intended to protect. Controls enforced in a separate hardware trust domain remain effective even when the operating system, runtime, or application code is compromised, which motivates identifying exactly where each security decision is made and who can alter it.

Hardware-based security creates a root of trust anchored in silicon, giving higher-layer defenses a foundation that host software cannot rewrite. Machine learning systems deployed in edge devices, embedded systems, and untrusted cloud infrastructure often depend on that anchor because the attacker may have physical access, co-tenant access, or administrative access to parts of the stack.

The trust boundary becomes especially important when systems must satisfy regulatory requirements. Healthcare ML systems handling protected health information under HIPAA must implement technical safeguards such as access control, audit controls, integrity protection, and transmission security. Systems processing personal data within GDPR's territorial and material scope must also demonstrate appropriate technical and organizational measures, including data-protection-by-design principles. Hardware security does not satisfy these requirements alone, but it

provides the boot, runtime, key-management, and device-identity guarantees that software controls often assume.

Four hardware security mechanisms divide that foundation by separating startup integrity, runtime isolation, key custody, and device identity. Secure boot verifies whether the system started from authentic code, whereas trusted execution environments (TEEs) isolate sensitive computation while it runs on an untrusted host. Because both mechanisms rely on cryptographic keys and device identity, hardware security modules (HSMs) generate and use keys without exposing them to ordinary memory, and physical unclonable functions (PUFs) derive a hardware-bound identity that is harder to copy than a software credential. Table 13.7 compares the four primitives after each has been examined. They do not map one-to-one onto the four deployment decisions of table 13.6, and they work best when composed across hardware, firmware, and software boundaries rather than selected as standalone products.

13.8.1.1 Trusted execution environments

A TEE is first a trust-boundary decision: which parts of the model, key material, input data, or attestation path must remain protected even if the surrounding host is compromised. A trusted execution environment (TEE) is the hardware-isolated processor region that implements that boundary. TEEs enforce confidentiality, integrity, and runtime isolation within a stated threat model. Their guarantee depends on the trusted hardware, firmware, and code inside the protected boundary.

For machine learning, a TEE can keep model parameters and sensitive inputs confidential during inference and protect execution from host-level tampering. It can, for example, prevent malicious software on the same device from extracting model parameters, or keep computations on facial or fingerprint data inside the protected boundary. Correct execution of an approved model says nothing about whether its predictions are accurate or safe.

Isolation of this kind appears at two scales. On consumer devices, a dedicated security coprocessor such as Apple's Secure Enclave keeps credentials and biometric features protected even if the application processor kernel is compromised (Apple 2024b). In the data center, the relevant form is the confidential virtual machine, which protects an entire guest, including its model weights and user inputs, from the hypervisor and host operators, and which the fleet admits into a job only after remote attestation (section 13.6.4).

Remote attestation supports the distributed trust decision: before components exchange sensitive data or model updates, a verifier checks fresh evidence about measured code and its hardware-protected boundary against an admission policy. Binding a communication key to that evidence connects the accepted identity to the ensuing channel. Attestation alone does not establish the correctness of every runtime operation. For ML, protection follows the computation, so inputs, weights, and activations remain covered only while their processing and storage stay within the chosen boundary.

TEE implementations combine execution isolation, protected memory, integrity mechanisms, and protected key storage in different ways. The isolation boundary may surround an application region, a secure processor mode, or an entire virtual machine. Access controls restrict which principals can read protected state, while memory encryption and authentication can extend protection beyond the processor boundary where the architecture supports them. Persistent storage needs its own protection when data leave protected memory. Within the execution boundary, computation uses plaintext directly. The deployment analysis must therefore trace where plaintext exists, which components can access it, and which integrity and freshness checks the platform actually supplies.

Platforms differ mainly in the size of the protected boundary. ARM TrustZone[13] splits a mobile or embedded processor into secure and normal worlds. Intel SGX isolates individual enclaves inside an application's address space on servers. Confidential virtual machines such as those built on AMD SEV-SNP or Intel TDX move the boundary out to a whole guest operating system, which lets unmodified ML serving stacks run inside it at the cost of a larger trusted code base.

Apple's system-on-chip (SoC) Secure Enclave shows these boundaries in silicon (figure 13.21). Its dedicated processor, true random number generator (TRNG), public key accelerator (PKA), and cryptographic engines operate separately from the application processor, which communicates through a mailbox interface. A memory protection engine encrypts and authenticates the enclave's external memory, while protected key paths support storage encryption (Apple 2024b). On consumer

13 **ARM TrustZone**: Introduced in 2004, TrustZone partitions an ARM processor into secure and normal "worlds" with hardware-enforced isolation. Its presence on a device does not by itself protect an ML deployment: the model, keys, and data path must actually be placed inside the secure world, which depends on whether the device vendor exposes secure-world applications beyond key storage.

devices, Apple's Face ID applies this architecture by binding its biometric templates and comparison logic to the Secure Enclave rather than treating them as ordinary application data (Apple 2024a). This example demonstrates isolation of designated security operations and state; it does not imply that arbitrary application model weights or activations reside inside the Secure Enclave, because their protection depends on the execution and data paths the platform supports.

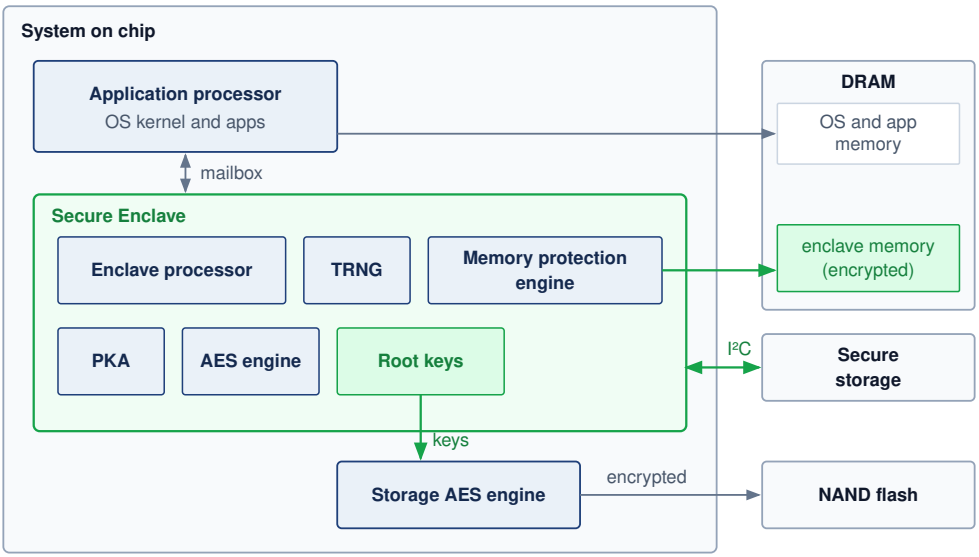

Figure 13.21: Secure Enclave Architecture: The Secure Enclave is an isolated region of the system on chip with its own processor, random-number generator, public-key accelerator, AES engine, and root keys. The application processor reaches it only through a mailbox, enclave memory in DRAM passes through a memory protection engine, and storage keys go directly to the inline AES engine, so a compromised application processor never holds them. Source: Apple.

Trusted execution environments carry costs in every C^3 dimension. Protected memory, memory encryption, and transitions between trusted and untrusted modes can add compute and latency, and battery-powered devices also pay in energy for secure memory accesses and cryptographic work. Development cost rises because code inside the boundary must follow isolation, attestation, and update protocols, and some deployments must also pass formal certification such as Common Criteria, an international standard for evaluating security products. Distributed training and federated inference add coordination work when their admission policy requires attestation and protected key exchange. A central verifier can perform admission checks; direct mutual attestation between every pair is not inherently required. TEEs therefore fit deployments whose threat model, whether adversarial users, cloud co-tenants, or malicious insiders, justifies that investment.

13.8.1.2 Protected-memory and overhead budgets

Hardware security features are budgeted like any other resource. Each primitive spends time at its boundary crossings and capacity in its protected memory, and both costs must be sized before a design commits to a primitive. Hardware security overhead must be measured for the chosen implementation. A secure-world transition has a cost per transition, not automatically a cost per inference: the result depends on how often the inference path crosses the boundary and what data it copies. Encryption, cache effects, and paging add separate costs. Energy overhead requires power and elapsed-time measurements; it cannot be inferred from a transition-latency figure alone.

Memory is the second budget. The protected working set competes for platform-specific memory capacity. Some implementations support protected paging beyond resident capacity, so overflow can impose substantial movement costs rather than an absolute model-size limit. Compression or partitioning changes those costs and may also change which values remain protected.

Security features scale differently than computational resources. Protected memory capacity, not host memory, sets how much of a model can remain resident inside a TEE without paging. On an older Intel SGX configuration with a 128 MB enclave page cache (the protected resident memory), ResNet-50's FP32 weights alone would fill about 76.2 percent of it. A ResNet-18 model is about 46.8 MB with FP32 weights or 11.7 MB with INT8 weights, while ResNet-50 is about 102.2 MB with FP32 weights or 25.6 MB with INT8 weights before activations and runtime buffers. The fleet's reference model sits far beyond that scale. The FP16 weights of the Archetype A (GPT-4/Llama-3) model occupy as many bytes as the 350 GB gradient it reduces each step, orders of magnitude more than an enclave page cache. Fleet-scale confidential computing therefore draws its boundary around a whole confidential virtual machine, an entire guest protected from the host, rather than an application enclave.

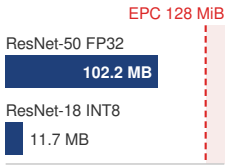

FP32 ResNet-50 weights alone nearly fill an older enclave page cache.

Beyond capacity, memory protection adds latency whose size depends on which threats are in scope. Encryption and authentication are distinct operations, and AES encryption alone does not verify integrity. An authenticated mode or a separate authentication mechanism is needed when tampering is in scope, and the latency must be measured for the actual data volume and hardware.

TEE assurance depends on its protected initialization and attestation path, not merely on when the host boots. Some enclave designs explicitly tolerate an untrusted host operating system, while other trusted environments depend on a verified platform startup chain. Secure boot complements these boundaries by authenticating the components included in that chain; the threat model must identify which firmware and initialization code remain trusted.

13.8.1.3 Secure boot

Secure boot is a mechanism that ensures a device only boots software components that are cryptographically verified under the platform's configured trust policy. At startup, each stage of the boot process, comprising the bootloader, kernel, and base operating system, is checked against a known-good digital signature. If any signature fails verification, the boot sequence is halted, blocking components that fail authentication. The chain rests on a trusted initial stage; signed but vulnerable software remains a risk.

A compromised boot process, particularly on embedded or edge hardware, can load malicious software before the ML runtime begins, enabling attackers to intercept model weights, tamper with training data, or reroute inference results. Such breaches can lead to incorrect or manipulated predictions, unauthorized data access, or device repurposing for botnets or crypto-mining.

For machine learning systems, secure boot authenticates the components included in its verification chain. Model artifacts receive this protection only when a verified loader checks their signatures or a signed manifest before loading them; boot verification alone does not authenticate weights read later from storage. Secure update flows must extend the same verification policy to replacement firmware, runtimes, and model artifacts.

Secure boot frequently operates in tandem with hardware-based TEEs to establish an unbroken chain of trust prior to code execution. The chain is a stepped verification handshake that pairs each trusted execution stage with an explicit cryptographic authentication gate (Regenscheid 2018), as figure 13.22 traces. Execution proceeds from an immutable boot root through a verified bootloader and operating system, halting and rejecting the component if any signature check fails. For machine learning pipelines, Stage 3 extends this chain beyond conventional system software: the verified OS or loader authenticates the model manifest and artifact digest before launching the authorized ML workload. The terminal stage emphasizes that while boot-time authentication blocks unauthorized or tampered binaries, dynamic runtime controls remain necessary to guard against post-boot attacks and data-dependent exploits.

Startup trust is what makes the runtime isolation described earlier meaningful, and it must survive updates. Apple's signed boot chain verifies the Secure Enclave firmware before any biometric data is processed and installs only software versions Apple is still signing, so a secure initial boot keeps its assumptions only through a controlled update path (Apple 2021b, 2024b, 2021c).

Secure boot also creates key-management and lifecycle obligations. Enterprises must protect signing keys, distribute trusted verification material, and support revocation or recovery when an authorized key is compromised. The affected trust chain depends on that key's authority and the platform's recovery design. Verification adds startup work, and every component covered by the

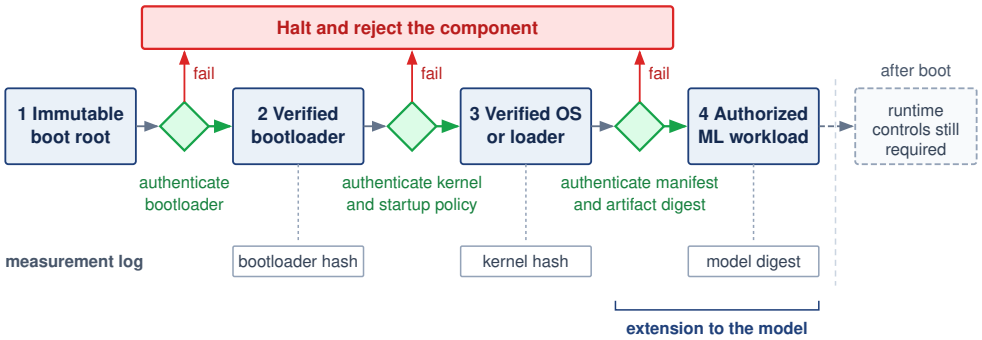

Figure 13.22: Secure Boot Sequence: Each trusted stage authenticates the next component before use. Extending the chain to a verified model loader permits model-artifact authentication; generic boot verification does not provide that extension automatically. Measured boot records evidence, whereas verification enforces the loading decision.

policy must remain authenticated through updates, whether by an individual signature or a digest bound into an authenticated manifest. The deployment must identify which components the chain actually covers.

At fleet scale, device identity and evidence of startup state support distinct admission decisions. The Device Identifier Composition Engine (Trusted Computing Group 2018) combines a device secret with software measurements to derive identities tied to the measured state. IEEE 802.1AR (IEEE 802.1 Working Group 2018) defines cryptographically bound device identities and their provisioning; a device identity alone does not establish which software is running. Fleet attestation (section 13.6.4) connects authenticated identity with acceptable state evidence before a node receives keys or joins a job.

Secure boot, when extended through the model loader and complemented by controlled updates, can establish that the configured authority approved the loaded components. Correct model behavior and freedom from vulnerabilities in the authenticated software remain separate claims. Runtime isolation and monitoring address different parts of that assurance argument.

13.8.1.4 Hardware security modules

While TEEs and secure boot provide runtime isolation and integrity verification, hardware security modules (HSMs) specialize in the cryptographic operations that underpin these protections. An HSM is a tamper-resistant physical device designed to perform cryptographic operations and securely manage digital keys. HSMs are standard in finance, defense, and cloud infrastructure; in an ML pipeline they protect the keys on which model integrity and data confidentiality depend.

HSMs provide an isolated, hardened environment for performing sensitive operations such as key generation, digital signing, encryption, and decryption. Unlike general-purpose processors, they are engineered to withstand physical tampering and side-channel attacks, and they typically include protected storage, cryptographic accelerators, and internal audit logging. HSMs may be implemented as standalone appliances, plug-in modules, or integrated chips embedded within broader systems.

In machine learning systems, HSMs can protect signing keys and the key-encryption keys that wrap, or encrypt, data-encryption keys. The distinction matters because a root key may remain inside the HSM while a workload uses a released data key inside its authorized execution boundary. Two distributed workflows illustrate this role. An HSM can protect the certificate authority's signing key used to certify enrolled devices; each device then proves possession of its own private key. Certification establishes identity, not benign behavior; an authenticated but compromised client can still submit poisoned updates. An HSM can also wrap a model-decryption key for a TEE's attestation-bound public key after the verifier accepts the evidence. Keeping weights out of plaintext host memory then depends on the TEE's key handling, loader, and execution path, not on the wrapping operation alone.

HSMs also play a role in maintaining the integrity of machine learning models. In many production pipelines, models must be signed before deployment to ensure that only verified versions are accepted into runtime environments. The signing keys used to authenticate models can be stored and managed within the HSM, providing cryptographic assurance that the deployed artifact is authentic and untampered. Similarly, secure firmware updates and configuration changes, regardless of whether they pertain to models, hyperparameters, or supporting infrastructure, can be validated using signatures produced by the HSM.

In addition to protecting inference workloads, HSMs can be used to secure model training. During training, data may originate from distributed and potentially untrusted sources. By holding the keys that sign training-data manifests and wrap data-encryption keys, HSMs can help ensure that training pipelines perform encryption, integrity checks, and access control enforcement securely and in compliance with organizational or legal requirements. In regulated industries such as healthcare and finance, such protections are often necessary to satisfy required safeguards. For instance, the HIPAA Security Rule requires technical safeguards such as access control, audit controls, integrity protections, and transmission security, while treating encryption and some integrity controls as addressable implementation specifications. GDPR lists pseudonymization and encryption as examples of appropriate technical and organizational measures where they fit the risk.

HSM deployment introduces a capacity and integration budget. Signing throughput, key-unwrapping latency, and network access to a remote HSM must be measured for the operations actually used, because the HSM sits on the signing and key-release path and its throughput, not the accelerator's, bounds how fast keys reach the fleet. The protected key boundary can also require changes to application interfaces and fleet provisioning. Certification and procurement requirements, including applicable cryptographic-module standards, constrain which devices and configurations can be used.

HSMs combine key custody with protected cryptographic operations. Resource-constrained embedded systems may need a smaller mechanism for a narrower task: deriving a device-specific secret without storing a persistent copy of that secret. Physical unclonable functions use a chip's physical properties for this purpose. They can support key derivation and identity, but do not replace the HSM's full cryptographic processing and policy-enforcement role.

13.8.1.5 Physical unclonable functions

PUFs[14] provide a hardware-intrinsic mechanism for cryptographic key generation and device authentication by exploiting physical randomness in semiconductor fabrication (Gassend et al. 2002). Unlike traditional keys stored in memory, a PUF generates secret values based on microscopic variations in a chip's physical properties, which are inherent to manufacturing processes and difficult to clone or predict, even by the manufacturer. Because of variation such as doping concentration and line edge roughness, chips fabricated from the same design masks differ measurably in timing or voltage behavior, and a PUF circuit amplifies those differences into a device-dependent response to each input challenge. An attacker may still predict responses from a vulnerable design without physically cloning the chip.

This challenge-response mechanism allows PUFs to serve several cryptographic purposes. They can be used to derive device-specific keys that never need to be stored externally, reducing the attack surface for key exfiltration. The same mechanism also supports secure authentication and attestation, where devices must prove their identity to trusted servers or hardware gateways. These properties make PUFs a natural fit for machine learning systems deployed in embedded and distributed environments.

For ML fleets, a PUF can reduce the need to inject and persistently store a unique secret in each device. The phone-class devices of Archetype C (Federated MobileNet) and the cameras and drones of edge deployments still need enrollment: the verifier must associate a device with trusted authentication evidence and manage its admission and revocation. A PUF-derived key can encrypt the on-device model at rest, so copying the stored file does not yield the device-bound key that decrypts it. Authentication can use enrolled challenge-response information or a credential backed by a derived key, depending on the design. Both protections have limits. Plaintext weights extracted during execution expose the shared model regardless of how the storage key was generated, and

[14] **PUF (physical unclonable function)**: Manufacturing variation can support device-specific key derivation without storing a persistent copy of the key. Separate device keys limit the reach of a storage-key compromise, but they do not prevent extraction of plaintext model weights during execution or attacks that predict a vulnerable PUF's responses.

successful authentication establishes which device is speaking, not that its runtime software is uncompromised.

Panel A of figure 13.23 states the common principle, and panels B to D show three PUF architectures that draw on different physical sources of variation (Gao et al. 2020). In each, two devices of the same design answer challenge C with different responses ($R_1 \neq R_2$), whether the response is an optical speckle pattern, an arbiter's response bit r set by racing signal paths, or the startup state of an SRAM cell. These responses are not perfectly repeatable under every operating condition, which motivates the error correction and calibration discussed next.

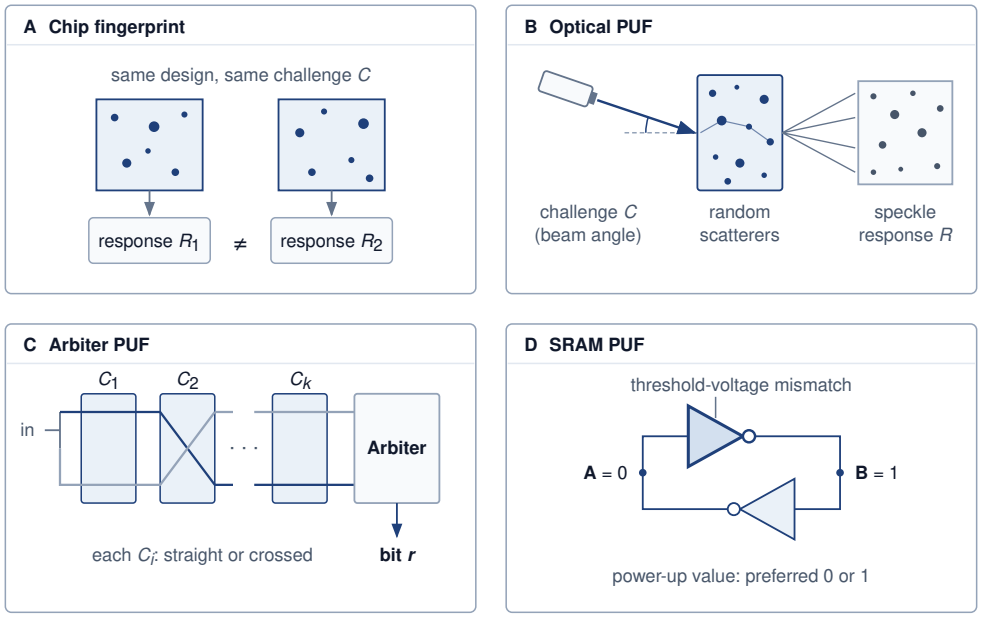

Figure 13.23: Physical Unclonable Functions: Each panel turns uncontrolled manufacturing variation into a device-specific response. Two chips of the same design answer the same challenge differently (A); a laser challenge scattered by a random token yields a speckle response (B); a challenge sets each arbiter stage straight or crossed, and the faster of two racing edges sets the response bit (C); and threshold-voltage mismatch decides the value a cross-coupled SRAM cell settles to at power-up (D).

PUFs also present engineering challenges. Their responses drift with temperature, voltage, and aging, so production designs add error correction or helper data and calibrate devices during fabrication, at some cost in layout area and test time. Managing large sets of challenge-response pairs raises questions of storage, consistency, and revocation, and the statistical structure of some PUF designs makes their responses predictable to machine learning-based modeling attacks if enough pairs are observed. Within those limits, PUFs reduce the key-management burden for large device fleets and supply a hardware-bound identity that complements secure boot, TEEs, and HSMs.

13.8.1.6 Mechanisms comparison

Trusted execution environments, secure boot, hardware security modules, and physical unclonable functions address distinct layers of the platform rather than substituting for one another. The design choice is which primitive owns which trust boundary, and table 13.7 sets the four side by side by function, use in ML, and trade-offs.

Table 13.7: Hardware Security Primitives Compared: Machine learning systems use diverse hardware defenses (trusted execution environments, secure boot, hardware security modules, and physical unclonable functions) to establish trust and protect sensitive data across the system stack. Each mechanism addresses a different security challenge, from runtime isolation and integrity verification to key management and device identity, and each carries its own trade-offs in performance and complexity.

Mechanism	Primary function	Common use in ML	Trade-offs
Trusted Execution Environment (TEE)	Isolated runtime environment for secure computation	Secure inference and on-device privacy for sensitive inputs and outputs	Added complexity, memory limits, performance cost; requires trusted code development
Secure Boot	Verified boot sequence and firmware validation	Authenticates boot components; model protection requires a verifying loader	Key management complexity, vendor lock-in; startup performance impact
Hardware Security Module (HSM)	Secure key generation, storage, and cryptographic processing	Signing ML models, securing training pipelines, verifying firmware	High cost, integration overhead, limited I/O; requires infrastructure-level provisioning
Physical Unclonable Function (PUF)	Hardware-bound identity and key derivation	Model binding, device authentication, protecting IP in embedded deployments	Environmental sensitivity, modeling attacks; needs error correction and calibration

Together, these hardware primitives support a defense-in-depth strategy for securing ML systems in adversarial environments. Their integration needs an explicit assurance argument: which components remain trusted, which evidence supports that trust, and which attacks fall outside the protection. Safety-sensitive devices and distributed learning systems need this argument as much as they need the mechanisms themselves.

Hardware primitives supply evidence and enforcement for selected platform properties. Deployment controls use those capabilities to decide which model artifacts may load and which callers may use them. The transition from platform assurance to application policy is therefore a composition problem: the loader and serving endpoint must enforce the intended permissions within the protected boundary.

13.8.2 Secure model deployment

Training a robust model inside a locked cluster accomplishes nothing if the exported artifact is packaged as an untrusted pickle file, stored unencrypted, or served behind an API key with no per-principal quota. A trusted platform establishes where a model runs; deployment controls decide who can load it, who can call it, and whether its behavior survives contact with users, operators, and attackers. As models move into edge devices, public APIs, and multi-tenant platforms, serialization, packaging, and endpoint access policies determine whether the integrity, confidentiality, and availability established during training survive inference.

Because section 13.5.1.4 already quantified the API-boundary query quotas, output perturbation, and pricing regimes that counter model extraction, deployment-stage controls here focus on artifact custody, cryptographic packaging, and endpoint authentication infrastructure, while the design-stage controls of section 13.8.4 shape the model's behavior and ownership evidence. Deployment controls also have to survive ordinary optimization work such as quantization, pruning, and knowledge distillation, where performance improvements must not compromise artifact integrity, ownership evidence, or access-control boundaries.

Once training is complete, the model must be securely packaged for deployment. Authenticated encryption protects stored artifacts when its keys remain inside the intended trust boundary; a signed manifest authenticates the expected artifact digest before loading. Obfuscation may raise reverse-engineering effort but does not provide the confidentiality contract of encryption. Runtime key custody and the loader's verification policy determine whether packaging protection survives execution.

Deployment environments must enforce access policies at inference endpoints. OAuth supplies delegated authorization; mutual TLS authenticates certificate-bearing peers; API keys identify credentials that the service maps to permissions and quotas. Role-based access control (RBAC) can restrict those permissions by operational role. These mechanisms support the query controls

in section 13.5.1.4, but a credential is not proof of a unique human or organization. An attacker who can obtain many accounts can still distribute queries across identities, so account issuance, cumulative budgets, and abuse detection remain part of the boundary.

An API key demonstrates possession of a credential and allows the backend to enforce its associated policies. Keeping keys out of source code reduces accidental exposure in version control, but environment variables alone do not prevent leaks through logs, process inspection, or a compromised runtime. Secret retrieval, scoped permissions, rotation, and redacted logging must be designed together. Model verification against an authenticated manifest, input limits, rate controls, and security-event logging complement credential management, while runtime monitoring (section 13.8.3) extends these protections into active operation.

13.8.3 Runtime system monitoring

A model binary that passed cryptographic signature verification at boot can still receive adversarial prompts, drift into unsafe generations, or have its weights modified in memory during operation. Secure deployment controls who can load the model and call the endpoint, while runtime monitoring checks whether the live system is still behaving like the approved system across cloud services, edge devices, and embedded hardware.

Runtime monitoring must decide whether an input should reach the model, whether the output is behaving normally, and whether the serving system is still the one originally deployed. The corresponding mechanisms are input validation, output monitoring, and system integrity checks.

13.8.3.1 Input validation

Input validation checks incoming data before it reaches a machine learning model. Format and range checks enforce explicit input contracts, while statistical and semantic checks estimate whether an input fits expected conditions. These checks can reject malformed inputs that disrupt preprocessing or inference and some adversarial inputs crafted to trigger incorrect predictions. Passing validation establishes only that the implemented checks accepted the input, not that the model will handle it correctly.

Small, carefully chosen changes to input data can cause models to make high-confidence but incorrect predictions (Goodfellow et al. 2015). Such inputs may satisfy ordinary type and range checks. Input validation therefore complements model robustness evaluation; it cannot replace testing how the learned function behaves on valid but adversarial inputs.

Validation techniques range from low-level checks of size, type, and value range to semantic filters, such as a face detector on an uploaded image or a speech activity detector on a voice clip. These checks reject inputs that violate the declared contract, but an adversarial input that satisfies the contract still reaches the model, and a semantic filter's own error rates belong in the validation assessment.

In generative systems, input validation can include prompt filtering against an application's content policy. A medical education application, for example, might reject requests to generate fabricated treatment claims. Filters may use keyword lists, regular expressions, or lightweight classifiers that estimate prompt intent. Their effectiveness depends on the policy, detector, and inputs tested, including paraphrases and adversarial prompts. They cannot establish that an accepted prompt will produce an acceptable output, which is why output monitoring remains necessary.

In some applications, distributional checks are also used to assess whether incoming data statistically resembles what the model saw during training. A computer vision pipeline, for instance, might compare the color histogram of an input image to a baseline distribution and flag outliers for manual review or rejection. These validations can be lightweight (heuristics or threshold rules) or learned (small models trained to detect distribution shift or adversarial artifacts). Their value depends on which invalid or hostile inputs they reject and how often they reject legitimate inputs. Since an accepted input can still induce an incorrect result, the next boundary is the model's output.

13.8.3.2 Output monitoring

Even when inputs pass validation, adversarial or unexpected behavior may still emerge at the model's output. Output monitoring observes confidence, prediction entropy, class distribution, or response patterns to flag deviations from expected behavior. A monitor can contain a detected

failure when its fallback acts before the output causes harm. Missed detections, delayed alerts, and incorrect fallbacks limit that protection, so detection quality and response timing belong in the same evaluation.

The signal a monitor watches depends on the model's modality. For a discriminative classifier, unexpected changes in confidence, class frequencies, or output entropy can trigger investigation, but do not identify their cause. A legitimate workload change, model defect, or attack can produce similar deviations. Content-moderation monitors can compare predicted "safe" labels with auxiliary reference signals; fraud-detection monitors can flag sudden score changes during high-volume periods. In each case, the system compares the live stream with an expected distribution and applies a chosen response, such as human review or a conservative baseline, while investigation distinguishes the possible causes.

Generative models, such as text-to-image systems, introduce unique output monitoring challenges. These models can produce high-fidelity imagery that may inadvertently violate content safety policies, platform guidelines, or user expectations. To mitigate these risks, postgeneration classifiers are commonly employed to assess generated content for objectionable characteristics such as violence, nudity, or brand misuse. These classifiers operate downstream of the generative model and can suppress, blur, or reject outputs based on predefined thresholds. Some systems also inspect internal representations (for example, attention maps or latent embeddings) to anticipate potential misuse before content is rendered.

Prompt filtering alone is nonetheless insufficient. Text-to-image systems can be manipulated through implicitly adversarial prompts, queries that appear benign but lead to policy-violating outputs, as the Adversarial Nibbler red-teaming project demonstrated on models such as Stable Diffusion (Quaye et al. 2024; Parrish et al. 2023). Such prompts bypass filters because the risk arises from model behavior during generation rather than from lexical cues in the prompt.

Prompts that pass a filter can still yield graphic outputs, such as the three blurred generations in figure 13.24, so an output needs assessment against the application's policy even when the prompt appears acceptable. Diagnosing how a prompt bypassed a filter requires prompt, model, and filter records alongside the outputs, and the output classifiers that catch these cases need their own errors evaluated.

Figure 13.24: Adversarial Prompt Evasion: Implicitly adversarial prompts bypass typical content filters by triggering unintended generations, revealing limitations of solely relying on pregeneration safety checks. These examples show the need for post-hoc content analysis as a complementary defense layer for robust generative AI systems. Source: (Quaye et al. 2024).

For language generation, output monitoring targets toxicity, hallucinated claims, and off-distribution responses. For example, a customer support chatbot may be monitored for keyword presence, tonal alignment, or semantic coherence. If a response contains profanity, unsupported assertions, or syntactically malformed text, the system may trigger a rephrasing, initiate a fallback to scripted templates, or halt the response altogether.

Output monitoring can combine rule-based heuristics with learned detectors trained on historical outputs. These detectors feed alerts into incident response pipelines or block outputs before release. Adversarial training aims to improve the model's behavior, while monitoring adds a separate opportunity to detect and contain failures during operation. Any claimed damage bound must account for missed detections, response latency, and the consequences of the selected fallback. Guardrail classifiers such as Llama Guard and ShieldGemma (section 13.5.5) occupy the same position in the runtime stack, so their own error rates belong in the deployment's safety evaluation.

13.8.3.3 Integrity checks

While input and output monitoring focus on model behavior, system integrity checks compare model files, execution environments, and serving infrastructure with approved state. Their coverage depends on what is measured, when it is measured, and whether the measurement path remains trusted. Verifying a file before loading detects covered artifact changes at that point; it does not establish that all subsequent execution remains untampered. One of the most common integrity mechanisms is cryptographic model verification: before a model is loaded into memory, the system computes a cryptographic hash (for example, SHA-256) of the model file and compares it against the expected digest in an authenticated, signed release manifest.

Access control and audit logging complement cryptographic checks. ML systems should restrict access to model files using role-based permissions and monitor file access patterns. For instance, repeated attempts to read model checkpoints from a nonstandard path, or inference requests from unauthorized IP ranges, may indicate tampering, privilege escalation, or insider threats. In cloud environments, container- or VM-based isolation[15] helps enforce process and memory boundaries, but these protections can erode over time because of misconfiguration or supply chain vulnerabilities.

For example, a regulated healthcare ML deployment can organize integrity evidence around three gates. The loader compares the model artifact's cryptographic hash with an authenticated, signed release manifest. Dependency validation checks approved package versions and authenticated artifact digests; a pinned version number names a package without proving that its contents are the approved ones. Platform attestation checks measured state against an approved configuration, with coverage determined by the attestation mechanism rather than an assumption of continuous execution verification. These gates supply evidence relevant to HIPAA[16] and other applicable obligations, but the organization must assess the complete system and its operational use to establish compliance.

Some systems also implement runtime memory verification, such as scanning for unexpected model parameter changes or checking that memory-mapped model weights remain unaltered during execution. While more common in high-assurance systems, such checks are becoming more feasible with the adoption of secure enclaves and trusted runtimes. Together, system integrity checks and input/output monitoring provide complementary evidence: integrity checks can identify changes to measured artifacts even when predictions appear normal, while behavioral monitors can flag failures even when every artifact matches its approved version. Their combined evidence guides investigation and containment without equating artifact authenticity with correct model behavior.

Integrity checks become ML supply-chain controls only when they connect every deployable artifact. A signed model file is insufficient if the training dataset, dependency lockfile, preprocessing code, registry promotion event, container image, and deployment attestation cannot be tied to the same release record. That release record makes the provenance bound on trust (principle 22) operational, and it gives rollback and forensic analysis a concrete object. Operators can ask which data snapshot produced the checkpoint, which dependencies built the serving image, which registry principal promoted it, and which production endpoints loaded it.

13.8.3.4 Response and rollback

A detected breach, anomaly, or performance degradation must lead to a defined response. The immediate task is containment; recovery then restores a trustworthy service and addresses the cause. ML systems add model checkpoints, training-data lineage, and learned behavior to the software versions, credentials, and runtime state that incident responders already need to examine.

Detection thresholds and automated rollback triggers work here as they do for any production model; security changes the signal set and the order of operations. Escalation can come from input validation (a rise in invalid inputs), output monitoring (a drop in prediction confidence), or integrity checks (a failed model signature verification). A security trigger must also preserve evidence before the automated response changes the system's state, because a rollback that overwrites the suspect artifact destroys the record an investigation needs.

One common strategy is model rollback, where the system reverts to a previously verified release. If a newly deployed fraud detector misclassifies transactions, an earlier checkpoint can restore service only if it remains compatible with the current data path and does not share the defect or compromise. Rollback therefore requires an authenticated release record covering weights,

[15] **Container vs. VM isolation**: Containers share the host kernel, which keeps overhead low and model scale-out fast but leaves a kernel attack surface shared with co-tenants. VMs separate guest kernels and memory at additional resource cost, yet hypervisor flaws and shared-hardware side channels can still enable cross-tenant attacks on a multi-tenant serving fleet.

[16] **HIPAA ML requirements**: For electronic protected health information, the HIPAA Security Rule treats encryption at rest and in transit as addressable specifications, which an organization implements when reasonable and appropriate or replaces with a documented equivalent (U.S. Department of Health and Human Services 2005). Its audit-control standard requires mechanisms that record and examine activity in systems holding that information, scoped by risk analysis rather than a mandate to log every read. Cloud services handling it generally require Business Associate Agreements. In an ML pipeline, these duties become risk-based audit coverage, protected checkpoints, controlled access, and evidence that the safeguards are operating.

↩ Automated rollback triggers for production models are established in section 12.5.6.2.

preprocessing, dependencies, and configuration, together with criteria for accepting the fallback under current conditions.

In high-availability environments, model isolation can help contain failures. Removing a suspect instance from load balancers stops new serving traffic, while restricting its network and credential access limits other actions. Preserving an isolated copy supports forensic analysis as unaffected replicas serve requests. The canary stages of section 12.5.5.2 serve a different purpose, evaluating a candidate release with limited live exposure. They do not substitute for quarantining a known compromised instance, because a small traffic share does not make continued use of a compromised model safe, and a rare trigger may never appear during the canary window.

Traffic throttling is another immediate response tool. If an adversarial actor is probing a public inference API at high volume, the system can rate-limit or temporarily block offending IP ranges while continuing to serve trusted clients. This containment technique helps prevent abuse without requiring full system shutdown.

Once immediate containment is in place, investigation and recovery can begin. Evidence may include protected input logs, parameter deltas between model versions, or memory snapshots from inference containers. Generative systems add prompts, retrieval context, guardrail decisions, and secret-access events. These records help reconstruct what instructions and data the model received, what it produced, and which controls intervened; a guardrail log records the detector's decision rather than proving that it caught every violation. Evidence collection must also protect the sensitive content it captures through scoped access, redaction, retention limits, and integrity controls. The aim is enough trustworthy context to distinguish privacy breaches, model-integrity failures, and ordinary model errors without creating an uncontrolled second repository of private data. Applicable incident-notification obligations remain part of the response process.

Recovery must repair the boundary that failed. A stolen credential requires revocation and controlled replacement; a compromised loader requires a trusted deployment path; poisoned training data may require dataset repair and retraining. Changes must pass through a secure update process using authenticated artifacts, trusted build pipelines, and validated data. For a prompt injection incident, recovery should examine how untrusted content gained authority over instructions or data access; additional training and filtering may help, but cannot replace enforcing the affected authorization boundary. Postincident review then follows the platform practice of section 12.7.4, with a security review adding three records to it: which boundary failed, which credentials and artifacts were revoked or rebuilt, and which attack cases now enter the evaluation suite.

13.8.4 Secure model design

Hardening the serving container does not stop an attacker from querying a high-confidence prediction on an out-of-distribution input or stripping ownership metadata from a stolen weight file. While deployment and runtime controls act on the platform around the model, the model tier above them acts on the learned function itself. Many of the vulnerabilities that platform controls try to catch at serving time can be reduced before deployment by choosing what the model can learn, how it behaves under uncertainty, and what evidence of ownership it carries.

One design strategy is to limit when the system acts on a model prediction. Confidence calibration makes reported confidence more meaningful under the evaluated conditions, while abstention permits the system to defer uncertain predictions. Neither mechanism establishes robustness by construction: an adversarial input may still receive high confidence. Output smoothing and other regularization methods likewise require evaluation under the specified attack. These design choices support behavioral robustness testing rather than replacing it.

Another design-stage consideration is the use of model watermarking,[17] a technique for embedding verifiable ownership signatures directly into the model's parameters or output behavior (Adi et al. 2018). A watermark might be implemented, for example, as a hidden response pattern triggered by specific inputs, or as a parameter-space perturbation that does not affect accuracy but is statistically identifiable. Abstention and watermarking therefore serve different purposes: abstention controls when the system acts, while watermarking supplies evidence about model origin, and each needs its own evaluation.

[17] **Model watermarking**: In trigger-set watermarking, the owner's secret inputs reliably identify the model while normal test accuracy is preserved (Adi et al. 2018). The systems challenge is robustness, because a watermark must survive fine-tuning, pruning, and distillation, the same operations an attacker uses to strip ownership marks from a stolen model.

Smaller or compressed models can reduce the protected working set and make an isolation boundary less costly to implement. That resource benefit buys no resistance to extraction, memorization, or side channels. A simpler model may be easier to approximate from queries, and leakage depends on the implementation and exposed outputs rather than parameter count alone.

In high-risk applications, designers may also prefer models whose decisions reviewers can inspect, which supports investigation of vulnerabilities and biases without establishing their absence. All of these model design choices reflect trade-offs between accuracy, robustness, transparency, and system complexity, and making them before deployment gives engineers an opportunity to evaluate the resulting behavior and resource cost together. Ordinary accuracy evaluation can still miss one risk entirely. A model can retain training examples and later expose them through its outputs, the leakage of section 13.5.2, which design choices at this tier do not remove.

↪ Where explanations belong in a deployed system, and what they cost, is developed in section 16.7.

13.8.5 Privacy-preserving data techniques

Addressing that training-data exposure moves the defense to the top tier of the stack, where the data layer forces an explicit decision about who may observe raw records, individual updates, and released results. Leaving raw typing logs on phones via federated learning can still leak private text through unmasked gradients, just as replacing patient records with a GAN-generated synthetic dataset can still leak records that the generator memorized. Keeping records centralized, leaving them on client devices, computing over encrypted values, and generating synthetic substitutes therefore change different parts of that exposure. That decision is architectural before it is algorithmic because it fixes the system variables tracked throughout the defense stack: information leakage, model utility, runtime cost, and the strength of the guarantee. ML workflows often depend on raw, high-fidelity data to train effective models, so privacy-preserving techniques should be read as different ways to move or bound that exposure rather than as interchangeable tools.

Each technique excludes a different observer. Federated learning keeps raw records away from the central server, secure aggregation hides individual updates from the coordinator, homomorphic encryption and secure multiparty computation hide inputs from the parties that compute on them, and synthetic data withholds the source records from the people who use the substitute. De-identification and k-anonymity, which strip or coarsen identifying fields, by contrast change only the shape of a released table.

13.8.5.1 Federated learning

A release guarantee and a data-location requirement answer different privacy questions. Differential privacy bounds the effect of a protected contribution on released results; federated learning keeps raw training records on participating clients. Policy, consent, or data-use agreements may forbid centralizing those records, making federated architectures a natural fit. The updates exchanged during training still require their own exposure analysis.

Federated learning restructures the learning process so that client data remain decentralized while model updates are communicated. Data locality alone does not establish privacy, however, because the gradient updates that clients transmit can leak sensitive information about the examples that produced them. Decentralization is necessary when the deployment contract requires local records, but a trusted curator applying differential privacy to centralized data can provide a guarantee under a different trust model.

↩ Federated coordination protocols, weighted averaging, and client update cycles are established in section 11.6.

13.8.5.2 Gradient leakage and privacy-preserving computation

Gradient inversion attacks (such as Deep Leakage from Gradients) demonstrate reconstruction of training examples from exposed gradients under the studied conditions (Zhu et al. 2019; Nasr et al. 2019). Recovery depends on the model, batch, update processing, and attacker's knowledge; it is not guaranteed for every client update. The gradient nevertheless provides an information channel even when raw data never leave the device. One defense is the secure aggregation protocol, which limits the server to an aggregate under specified participation and collusion assumptions.

Each client's update carries pairwise masks that cancel in the server's sum, so the coordinator recovers the aggregate without receiving any unmasked update. That protection holds only under stated conditions. Enrollment must be authenticated, the participating clients must complete the round, and at least two noncolluding clients must keep their shared masks from the coordinator.

↩ Secure aggregation and its pairwise masks, which reveal only the sum of client updates to the coordinator, are established in section 11.7.4.

A production protocol must therefore formalize its collusion threshold and its dropout recovery, because a client that departs mid-round strands masks that no longer cancel. The aggregate itself still discloses information, potentially including an individual contribution when combined with other known contributions, so hiding the messages does not by itself provide differential privacy.

Gradient leakage motivates both secure aggregation protocols (Bonawitz et al. 2017) and differential privacy, but the trust boundary determines which version of DP applies. Central differential privacy assumes a trusted curator or aggregation service can see raw records or updates before releasing a randomized result. Local differential privacy randomizes each client's protected contribution before the server sees it, reducing trust in the server but often paying a larger utility cost. Secure aggregation limits the server to an aggregate rather than individual updates; a separate, accounted DP mechanism can bound what the released result reveals. Noise added after the server sees the exact aggregate protects the release audience, not that server. Protecting against an untrusted coordinator requires randomization or a trusted computation boundary before it observes the protected result. Compression can change the information in an update, but communication savings alone establish no privacy bound. Each of these protections also carries a C^3 price. Masking adds client compute, key agreement and masked uploads add communication, and dropout recovery adds protocol-specific coordination governed by participation thresholds and deadlines.

↩ Gradient compression as a lever against bandwidth scarcity is established in section 5.9.

Lighthouse 13.1: Archetype C (Federated MobileNet): Google Gboard federated learning

For the Archetype C mobile/federated workload, Google's Gboard keyboard (whose on-device next-word suggestion and autocomplete inference interfaces appear in figure 11.1) provides a concrete production baseline for learning from data retained on edge devices. The original keyboard study demonstrates federated training, and the systems report describes the production orchestration challenges (Hard et al. 2018; Bonawitz et al. 2019). Those reports do not by themselves establish that every described model combines secure aggregation with a formal DP guarantee. The engineering pattern separates three contracts: FL keeps raw typing data local, secure aggregation can hide individual messages under its protocol assumptions, and an accounted DP mechanism can bound what the released model reveals. A deployment claim must identify which contracts its actual training run satisfies. For this archetype the threat model therefore starts after the raw data remain on the device, with what each device's update reveals and who can observe it.

Federated learning changes where raw data reside, while differential privacy bounds the effect of a protected contribution on released results. Cryptographic computation addresses a third boundary: who can observe inputs and intermediate values during execution. Homomorphic encryption (HE)[18] and secure multiparty computation (SMPC) support computation over inputs represented as ciphertexts or protocol-specific shares (Gentry 2009; Yao 1982). Their costs depend on the chosen protocol and representation. Many HE schemes pack multiple values into one ciphertext so that each operation acts on all of them, and rotations move values between those packed positions. Quantization can help some encrypted computations, but its effect must be assessed together with packing, rotations, nonlinear-function approximations, and required accuracy. Hardware acceleration changes execution cost without removing those mathematical constraints.

[18] **Homomorphic encryption**: Operations on ciphertexts correspond to operations on the underlying plaintexts. Gentry's 2009 work supplied the first fully homomorphic construction; it did not make arbitrary ML workloads immediately practical. Inference cost depends on circuit depth, packing, rotations, numerical representation, and the chosen security parameters.

In the case of HE, evaluating a supported function on a ciphertext produces a ciphertext that decrypts to the corresponding plaintext result:

$$c \leftarrow \text{Enc}(pk, m), \qquad \text{Dec}(sk, \text{Eval}(pk, f, c)) = f(m)$$

This property supports inference over encrypted inputs in an untrusted computing environment. Its confidentiality guarantee depends on the scheme and key custody; it does not by itself prove that the server evaluated the intended function correctly. The deployment must also measure the cost of representing that function as supported encrypted operations, including any approximation that changes prediction quality.

Napkin Math 13.2: The cost of trusted compute

Problem: A team deploying a health-monitoring model compares three security levels:

1. **Plaintext**: Standard inference.
2. **Encrypted transport (AES)**: Model and data are encrypted at rest and in transit.
3. **Encrypted compute (fully homomorphic encryption [FHE])**: Inference performed on encrypted data.

Treat the following overheads as assumed scenario inputs, not universal AES or FHE performance ratios. The comparison is whether these configurations satisfy the application's latency budget.

Math: Inference latency scales by the complexity of the security protocol.

1. **Plaintext**: 20 ms.
2. **AES-256**: 20 ms + 0.5 ms = 20.5 ms (negligible overhead).
3. **FHE**: 20 ms $\times$ 10,000× = 200 seconds.

Systems insight: Under the assumed overheads, protected computation takes longer than encryption around plaintext execution. Whether FHE fits an application depends on its latency budget and the measured encrypted circuit; the multiplier here does not generalize across models or implementations. Hardware isolation provides a different trust boundary in which computation uses protected plaintext. Comparing that option requires its own measured overhead and an explicit hardware trust assumption.

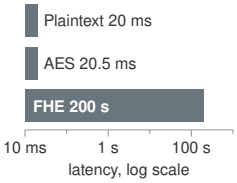

Under the assumed overheads, AES adds almost nothing; FHE adds four orders of magnitude.

Homomorphic computation costs depend on the supported operations, circuit depth, numerical representation, and chosen security parameters. Supporting arbitrary computation does not by itself determine the cost of a particular inference circuit. Feasibility requires measuring the protected workload against its latency, throughput, and utility requirements. The hardware-enforced isolation of section 13.8.1.1 offers a different trade-off: plaintext computation inside a trusted boundary, with security dependent on that boundary's hardware and software assumptions.

SMPC,[19] by contrast, distributes computation across parties while protecting their private inputs under the protocol's stated corruption and collusion assumptions. In secret-sharing protocols, for example, each input is split into random shares that reveal nothing on their own, each party computes on the shares it holds, and only the recombined shares disclose the result. The protocol specifies which parties receive the output; privacy hides additional information beyond what those outputs and a party's own inputs reveal. This is particularly useful in joint training across institutions with strict data-use policies, such as hospitals or banks.

13.8.5.3 De-identification and k-anonymity

Removing names and account numbers is the oldest privacy technique for a released dataset, and on its own it fails. The fields that remain act as quasi-identifiers, which an attacker can join against auxiliary data to single out a person. Linkage of the Netflix Prize dataset[20] showed the failure at scale. Researchers inferred individual movie preferences from records released without names by matching them to public ratings, not by querying a trained model (Narayanan and Shmatikov 2006). The Strava heatmap of section 13.3 showed the aggregate form of the same failure, in which sparse regions singled out the few people behind them.

A release satisfies **k-anonymity**[21] when every combination of its designated quasi-identifiers occurs in at least k_{anon} records, so no record can be singled out from fewer than k_{anon} candidates. The property belongs to a released table, not to a single query, so transforming one request does not establish it. It also constrains only the shape of the released table. When every record in a group shares the same sensitive value, or the attacker holds auxiliary knowledge about a person, knowing which group holds that person still discloses the attribute. The property therefore prevents an observer of the release only from singling out a record, and it bounds nothing about what a model

19 **SMPC performance overhead**: Cryptographic computation adds protocol-specific communication and local work. The bottleneck depends on round count, network latency, arithmetic representation, and the adversary model; a single slowdown factor does not characterize all deployments. Moving some layers outside protection changes the information exposed and therefore requires a new privacy analysis.

20 **Netflix deanonymization**: As few as 8 movie ratings with approximate dates uniquely identified 99 percent of records in the anonymized Netflix Prize dataset, and correlation with public IMDb ratings re-identified users (Narayanan and Shmatikov 2008). For ML training data, removing direct identifiers does not establish privacy when retained records can be linked to auxiliary information.

21 **k-anonymity**: A property of a released dataset (Sweeney 2002). Generalization, such as coarsening dates or postal codes, and suppression of rare rows can achieve the property at the cost of information loss. It does not bound inference from homogeneous sensitive attributes or auxiliary knowledge and is not a per-request privacy guarantee.

trained on the data reveals. Differential privacy (section 13.9) closes that gap by bounding how much one record changes what is released rather than how the released table looks.

13.8.5.4 Synthetic data generation

Synthetic data generation changes the privacy decision from hiding raw records to replacing them with generated surrogates. Generated data can let downstream researchers train models without receiving the source records, but the generator and its outputs must still be evaluated for leakage.

Synthetic data generation works by training a generative model (such as a GAN, variational autoencoder (VAE), or diffusion model) on the original sensitive dataset, then using this trained generator to produce new artificial samples. The intended result preserves useful population patterns without reproducing individual records. Learning a distribution does not ensure that separation, however; a generator can memorize its training data, so synthetic sampling alone cannot establish privacy.

Synthetic data generation typically advances through three ordered stages. First, in distribution learning, a generative model G_θ trains on real data $\mathcal{S}_{\text{real}} = \{x_1, x_2, \ldots, x_n\}$ to approximate the underlying data distribution $p(x)$. Next, synthetic sampling produces new records $\mathcal{S}_{\text{synthetic}} = \{G_\theta(\mathbf{z}_1), G_\theta(\mathbf{z}_2), \ldots, G_\theta(\mathbf{z}_m)\}$ by sampling latent vectors from random noise $\mathbf{z}_i \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. Finally, empirical validation measures selected aspects of fidelity to $\mathcal{S}_{\text{real}}$ and tests for memorization or linkage under the chosen attack model. Sampling can preserve useful statistical properties without distributing the source table (Goncalves et al. 2020), but passing those tests does not prove that every training record is protected. A formal release guarantee requires an additional mechanism, and validation that accesses sensitive source data must itself remain within the intended privacy boundary.

While appealing, synthetic data generation faces important limitations. Generative models can suffer from mode collapse, failing to capture rare but important patterns in the original data. Adversaries can also extract information about the original training data through generative model inversion or membership inference attacks. The privacy protection depends heavily on the generative model architecture, training procedure, and hyperparameter choices, making it difficult to provide formal privacy guarantees without additional mechanisms like differential privacy.

Consider a hospital that wants to share synthetic data for ML research. It trains a GAN on 10,000 patient records containing demographics, lab results, and diagnoses, seeking to preserve useful relationships among those features. The hospital applies differential privacy during training and uses a tested accountant, specifying whether the protected contribution is one record or all records from one patient. Before release, it evaluates clinical fidelity and subgroup utility, since an acceptable overall score can hide poor performance for rare conditions. Memorization tests supply additional empirical evidence, while the formal guarantee comes from the accounted mechanism. DP limits the effect of a protected contribution; it does not prevent learning population patterns or all sensitive inferences based on them.

Across federated learning, differential privacy, homomorphic encryption, secure multiparty computation, and synthetic data, the privacy-preserving approaches differ in the guarantees they offer and in their system-level implications. Production designs often combine them, for example by applying differential privacy inside a federated setup or by reserving homomorphic encryption for the few inference stages whose inputs justify its cost. The choice of mechanism therefore depends on computational constraints, deployment architecture, and regulatory requirements. Once the formal machinery of differential privacy is in place, these methods can be weighed on privacy guarantee, runtime overhead, operational maturity, and use case fit together rather than as independent properties (table 13.8).

The defense stack now covers each tier, from hardware roots of trust through deployment, runtime, and model-design controls to the data path, where federation, secure aggregation, encrypted computation, and synthetic data each move raw records or individual updates away from a particular observer. Data locality, message hiding, encrypted computation, and synthesis alone do not bound what the released model reveals. A secure training cluster and confidential links can still publish a model that exposes memorized examples, as the GPT-2 case showed, because the leak travels through an authorized output rather than an unauthorized access. Closing that gap requires a

guarantee stated about the release itself, with the protected contribution, the randomization, and the accumulated loss across releases all made explicit.

13.9 Differential Privacy

Suppose an adversary queries a medical diagnosis model using the exact attributes of a known patient and compares the answers with what a model trained without that patient would return. Encryption and access control cannot help here. The adversary is an authorized caller, and the query reaches only the trained model, never the stored records. A difference in one pair of predictions alone does not prove a privacy violation. The relevant question is how much the distribution of released outcomes permits an adversary to distinguish datasets with and without that patient, relative to a stated privacy bound.

Differential privacy constrains a randomized training or analysis mechanism so that its output distributions remain close on adjacent datasets. The adjacency definition specifies the protected contribution, such as one record or all records from one person. The information leakage invariant (principle 23) holds that any observable output can carry information about a protected training record, so a privacy claim must name the protected unit and the observations available to an adversary. Differential privacy answers that requirement by bounding the additional evidence a release supplies about the named contribution. For a fixed mechanism and workload, tighter privacy bounds require more noise or fewer data-dependent releases, and the first cost appears as lost utility. Recovering that utility with more data or more training steps raises the training compute, the T_{compute} term of the Fleet Law, and per-example gradient clipping in differentially private stochastic gradient descent (DP-SGD) lowers throughput on its own. Whether a tighter bound costs accuracy, compute, or both is therefore a design choice, which places DP in the deployment architecture rather than only in a legal checklist.

Definition 13.5: Differential privacy

Differential privacy is a mathematical guarantee that bounds changes in the output distribution of a randomized algorithm under a specified change to one protected contribution. The neighboring relation defines that contribution and the permitted change. Approximate differential privacy permits a multiplicative bound and additive slack, formalized by the (ϵ, δ) bound $\Pr[\mathcal{M}(X) \in \mathcal{S}] \leq e^{\epsilon} \Pr[\mathcal{M}(X') \in \mathcal{S}] + \delta$ for adjacent datasets X and X' and any measurable output set $\mathcal{S}$ (Dwork et al. 2006; Dwork and Roth 2014).

1. **Significance**: The privacy budget ϵ converts an algorithmic property into an engineering currency that can be allocated, composed, and depleted. Production reports and benchmark studies show that privacy budgets are deployment-specific: Apple describes local-DP telemetry deployments with per-event ϵ values chosen by use case, while DP-SGD studies show accuracy and compute trade-offs across benchmark tasks (Apple Differential Privacy Team 2017; Abadi et al. 2016; Bu et al. 2020; De et al. 2022). Smaller ϵ raises the noise floor and degrades classification accuracy along the empirical frontier plotted in figure 13.26, where CIFAR-10 suffers double-digit accuracy drops even at $\epsilon = 8$. DP-SGD can also reduce training throughput because per-sample gradient clipping works against the batch-level parallelism that makes accelerator training efficient.
2. **Distinction**: Unlike de-identification, which removes identifiers from records and fails against re-identification through auxiliary data, and unlike statistical disclosure control, which depends on the attacker's model, differential privacy is a property of the algorithm, not of the data. Its guarantee holds against any adversary, including one with arbitrary side information about other records in the dataset.
3. **Common pitfall**: A frequent misconception is that an ϵ label applies independently to every data access. Sequential composition of k DP mechanisms, each with privacy loss ϵ_0, yields total loss at most $k\epsilon_0$ under basic composition, with tighter bounds available from mechanism-specific accountants (section 13.9.2). A system that runs DP-SGD steps or answers interactive DP queries against sensitive data must therefore track cumulative

loss and stop or redesign the release when the allocated budget is exhausted. Rotating cryptographic keys does not reset privacy loss, while querying a fixed DP-trained model is post-processing and incurs no additional training privacy loss.

A randomized mechanism $\mathcal{M}$ run on adjacent datasets X and X' produces two output densities that overlap without coinciding (figure 13.25, blue and orange curves). Overlap means that either dataset can produce outcomes in the same region; it does not make the distributions identical. For a discrete outcome o in Range$(\mathcal{M})$, relative likelihood can be expressed as $\Pr[\mathcal{M}(X) = o]/\Pr[\mathcal{M}(X') = o]$ when the denominator is positive. Continuous curves instead depict densities, whose ratio is not their vertical difference. The (ϵ, δ) guarantee applies to probabilities of measurable output sets: it bounds one probability by e^{ϵ} times the other plus δ. For a continuous output, those probabilities are areas under the curves over the selected set. The definition therefore constrains distributional change across all such sets, rather than declaring an overlapping region private or treating additive slack as a pointwise density-ratio allowance.

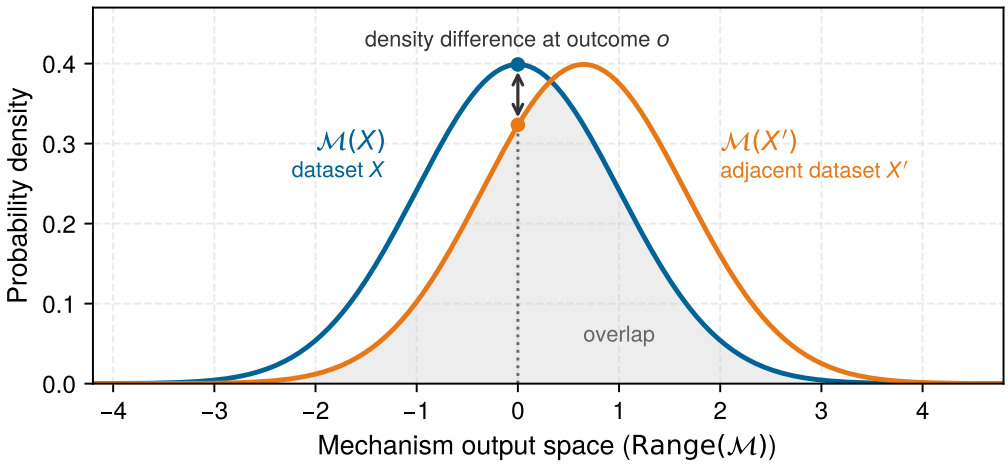

Figure 13.25: Differential Privacy and Output Distributions: These illustrative densities show how outputs can overlap without having identical likelihoods. Differential privacy bounds how much the output distribution can change between adjacent datasets X and X'. The privacy parameters limit, rather than eliminate, the information an observer can gain about any one individual's inclusion.

A randomized algorithm $\mathcal{M}$ is said to be ϵ-differentially private if, for all adjacent datasets X and X' differing in one record, and for all measurable output sets $\mathcal{S} \subseteq$ Range$(\mathcal{M})$, the following holds (Dwork et al. 2006; Dwork and Roth 2014):

$$\Pr[\mathcal{M}(X) \in \mathcal{S}] \leq e^{\epsilon} \Pr[\mathcal{M}(X') \in \mathcal{S}]$$

The parameter ϵ quantifies the privacy budget. In the pure-DP definition above, smaller values of ϵ impose a tighter likelihood-ratio bound; achieving that bound may require more noise or a different mechanism. No universal ϵ value denotes "strong," "moderate," or acceptable privacy: interpretation depends on adjacency, the additive slack δ in approximate DP, composition, and the consequences of inference. Under pure DP, the multiplicative factor e^{ϵ} bounds ratios of probabilities for the same output event on adjacent datasets. Approximate DP relaxes that event-probability inequality with additive slack, as developed below.

This bound limits how distinguishable the output distributions can be when a protected contribution changes; it does not make them identical or prevent inference of population-level patterns. Common DP mechanisms add calibrated noise to model updates or query responses, using distributions such as Laplace or Gaussian. Training techniques like DP-SGD[22] combine contribution bounds with calibrated noise and composition accounting. Their guarantee applies to the specified contribution and release, which must be identified before selecting the mechanism.

22 **DP-SGD (differentially private stochastic gradient descent)**: Introduced by Abadi et al. (2016), DP-SGD clips each per-sample gradient and adds calibrated Gaussian noise. Per-example clipping adds memory, recomputation, and reduction work that ordinary minibatch training avoids; vectorized clipping implementations recover much of the parallelism, so the throughput cost must be measured for the chosen model and implementation.

13.9.1 Mathematical foundations and privacy parameters

Attaching an ϵ label to a model without stating whether adjacency protects one record or one person, how many training steps composed, and what δ slack was allowed makes the privacy claim meaningless. Repeated training steps, retraining runs, or queries that access protected data may compose, while post-processing a fixed DP release incurs no additional privacy loss for that data. The mathematical foundations specify the guarantee, noise calibration, and accounting for those data-dependent operations so teams can define a valid accounting boundary before attaching an ϵ value to a deployed release.

13.9.1.1 The (ϵ, δ)-differential privacy formulation

The pure ϵ-DP definition provides strong guarantees but can be overly restrictive for practical ML applications. The relaxed (ϵ, δ)-differential privacy of the definition callout adds slack δ to the probability bound controlled by ϵ, and it is the starting point for noise calibration. Approximate DP is widely used with Gaussian perturbations in deep learning. Its role follows from the probability bound and noise calibration, not from unbounded support alone; Laplace noise also has unbounded support.

Formally, a randomized algorithm $\mathcal{M}$ satisfies (ϵ, δ)-differential privacy if for all adjacent datasets X, X' and all measurable output sets $\mathcal{S}$:

$$\Pr[\mathcal{M}(X) \in \mathcal{S}] \leq e^{\epsilon} \Pr[\mathcal{M}(X') \in \mathcal{S}] + \delta$$

The parameter δ is additive slack in this event-probability inequality, not generally the probability of a catastrophic disclosure. In practice, δ is commonly chosen smaller than the inverse dataset size, often $\delta < 1/D$ or smaller depending on policy and threat model; smaller δ tightens the guarantee at additional utility cost. For a dataset with 1 million records, $\delta = 10^{-12}$ specifies a small additive slack, but does not by itself assign a probability to every possible privacy failure.

13.9.1.2 Noise mechanisms for achieving differential privacy

ML systems implement differential privacy through carefully calibrated noise injection. The two primary mechanisms used in ML systems differ in their noise distributions and applicability. The first is the Laplace mechanism. For pure ϵ-DP, it adds noise drawn from $\text{Lap}(\Delta f/\epsilon)$, where Δf is the global sensitivity of the function (the maximum change in output when one record changes). For a query f, the privatized output is:

$$\tilde{f}(X) = f(X) + \text{Lap}\left(\frac{\Delta f}{\epsilon}\right)$$

The Laplace distribution has density $p(x) = \frac{\epsilon}{2\Delta f} \exp\left(-\frac{\epsilon|x|}{\Delta f}\right)$, with scale parameter $b = \Delta f/\epsilon$. For a fixed query and adjacency relation, tightening the bound from $\epsilon = 1.0$ to $\epsilon = 0.1$ increases the noise scale by $10\times$. Individual noise draws remain random; the scale controls their distribution rather than fixing their magnitude.

The Laplace derivation shows why the noise scale must match sensitivity. To prove that the Laplace mechanism achieves ϵ-DP, consider the ratio of output densities on adjacent datasets X and X'. Let $f(X) = v$ and $f(X') = v'$, where $|v - v'| \leq \Delta f$ by the sensitivity bound. For any output o, the density ratio is:

$$\frac{p(o|X)}{p(o|X')} = \frac{\exp(-\epsilon|o - v|/\Delta f)}{\exp(-\epsilon|o - v'|/\Delta f)} = \exp\left(\frac{\epsilon(|o - v'| - |o - v|)}{\Delta f}\right)$$

Using the triangle inequality, $|o - v'| - |o - v| \leq |v - v'| \leq \Delta f$. Therefore:

$$\frac{p(o|X)}{p(o|X')} \leq \exp\left(\frac{\epsilon \cdot \Delta f}{\Delta f}\right) = e^{\epsilon}$$

This establishes that the Laplace mechanism satisfies ϵ-differential privacy. The derivation reveals why sensitivity calibration is necessary, because the noise scale must match the maximum possible change in the query output to mask individual contributions.

Napkin Math 13.3: The cost of differential privacy

Problem: Consider computing the average salary of 1000 employees while guaranteeing privacy budget $\epsilon = 1$. The salaries range from \$0 to \$200,000. Assume a fixed, public employee roster and neighboring datasets that differ in one employee's salary, with both values inside these bounds. How much noise must the mechanism add?

Math:

1. **Sensitivity** (Δf): The maximum one person can change the sum is \$200,000.
2. **Privacy Budget** (ϵ): 1.
3. **Laplace Noise Scale** (b): the Laplace mechanism sets its noise scale to the sensitivity divided by the privacy budget, so $b = \Delta f/\epsilon =$ \$200,000 / 1 = \$200,000.
4. Impact on Mean: The mean-estimation noise comes from noise added to the sum with magnitude approximately \$200,000.
    - Error in the mean = \$200,000 / 1000 = \$200. The mechanism draws noise once, for the sum; dividing by the group size spreads that single draw across the mean.

Systems insight: Under this fixed-roster contract, the published mean receives zero-mean Laplace noise with expected absolute magnitude \$200; individual draws can be smaller or larger. For a dataset of $D = 100$, the expected absolute error grows to \$2,000. The noise needed to hide one person is fixed by the sensitivity, while the signal grows with the number of people, so small groups pay the largest relative cost and large datasets dilute it through the factor $1/D$.

The salary-average notebook above uses the central model of differential privacy. Each person supplies a bounded salary to a trusted analyst; the analyst computes the sum, adds one Laplace draw calibrated to the sum's global sensitivity, and divides by the fixed, public group size. The example protects a salary replacement within a public roster. If participation itself is private, adjacency must be add/remove, meaning adjacent datasets differ by the presence or absence of one person's records rather than by a change to one value. Under that relation the count becomes data dependent and needs its own privacy treatment, so division by that count is not automatically safe post-processing. This differs from local differential privacy, in which each person randomizes a value before submission and the analyst aggregates the randomized reports. Both models compare output distributions on adjacent inputs, but their mechanisms, trust assumptions, and utility are not interchangeable.

The Gaussian mechanism serves the ML case where the relaxed (ϵ, δ) guarantee is acceptable. It adds noise with standard deviation σ_{mech} calibrated based on ϵ, δ, and the ℓ_2 sensitivity $\Delta_2 f$:

$$\tilde{f}(X) = f(X) + \mathcal{N}\left(\mathbf{0}, \sigma^2_{\text{mech}}\mathbf{I}\right)$$

Unlike the Laplace mechanism, which achieves pure ϵ-DP, the Gaussian mechanism generally requires the relaxed (ϵ, δ)-DP formulation because shifted Gaussian densities can have arbitrarily large likelihood ratios in the tails. Bounding those tail events gives a sufficient way to establish the δ slack. The derivation works with the privacy loss, the log of the ratio between the two output densities at the observed output. Because the output is random, so is this log ratio, which is why it is called the privacy loss random variable.

For adjacent datasets with ℓ_2 sensitivity $\Delta_2 f$, the privacy loss at output $\mathbf{y}$ follows:

$$\mathcal{L}_{\text{priv}}(\mathbf{y}) = \ln \frac{p(\mathbf{y}|X)}{p(\mathbf{y}|X')} = \frac{\|\mathbf{y} - f(X')\|_2^2 - \|\mathbf{y} - f(X)\|_2^2}{2\sigma^2_{\text{mech}}}$$

When $\mathbf{y} = f(X) + \mathbf{z}$ for noise $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \sigma^2_{\text{mech}}\mathbf{I})$, expanding the difference of squared norms cancels the quadratic term $\|\mathbf{z}\|_2^2$ and leaves a constant shift plus a linear inner product between $\mathbf{z}$ and the dataset difference $f(X) - f(X')$. Unlike the Laplace log ratio, which the triangle inequality caps at a constant across the entire output space, this linear function of $\mathbf{z}$ grows without bound along the direction of the dataset difference. For an adjacent pair attaining the sensitivity bound, the privacy

loss is therefore a shifted Gaussian random variable with mean $\frac{\Delta_2 f^2}{2\sigma^2_{\text{mech}}}$ and variance $\frac{\Delta_2 f^2}{\sigma^2_{\text{mech}}}$, and δ bounds the probability that a noise draw lands far enough in the tail to push $\mathcal{L}_{\text{priv}}(\mathbf{y})$ above ϵ. For one application of the mechanism and $0 < \epsilon < 1$, a classic sufficient calibration is:

$$\sigma_{\text{mech}} \geq \frac{\Delta_2 f}{\epsilon}\sqrt{2\ln\left(\frac{1.25}{\delta}\right)}$$

This sufficient bound is a one-shot Gaussian-mechanism result within that parameter domain. Neither it nor the one-shot Laplace ratio $\Delta f/\epsilon$ of the salary example can set a DP-SGD noise multiplier, the ratio of the noise standard deviation to the clipping norm, by substituting the final training budget. DP-SGD clips sampled per-example gradients, adds Gaussian noise to their sum, and composes many such steps, so its budget requires a sampled-Gaussian accountant. Smaller ϵ or δ generally requires more noise, but the resulting accuracy change must be measured for the specific task and training procedure.

Napkin Math 13.4: One-shot Gaussian mechanism

Problem: Determine a sufficient noise standard deviation σ_{mech} for the Gaussian mechanism to achieve (ϵ, δ)-differential privacy for a single vector-valued query.

Variables: ℓ_2 sensitivity $\Delta_2 f = 1$, target $\epsilon = 0.5$, $\delta = 10^{-5}$.

Math: Classic sufficient bound requires $\sigma_{\text{mech}} \geq 9.7$.

Systems insight: This calculation applies to a single Gaussian release. DP-SGD repeats a Gaussian release at every optimization step, so its budget must be composed across thousands of steps, and the accountants of section 13.9.2.1 bound that composed total more tightly than adding up the per-step bounds would.

Gaussian noise fits the clipped-gradient mechanism because clipping explicitly bounds gradient norms in ℓ_2 space, making sensitivity analysis tractable. Rényi differential privacy accounting can retain mechanism-specific information to improve on generic composition bounds; that advantage belongs to mechanisms like clipped-gradient training and does not make Gaussian noise the right choice for every private computation.

13.9.2 Privacy budget composition

Releasing 10 models trained on the same patient cohort at $\epsilon = 1$ each does not maintain $\epsilon = 1$ protection; under basic composition the cumulative bound reaches $\epsilon = 10$. Because a full training run or a sequence of model releases applies a mechanism many times, privacy loss accumulates across every additional data-dependent DP release about the same protected population. Basic composition bounds the total by adding the per-release bounds, which is valid but often loose, while mechanism-specific accountants retain more information about each mechanism to obtain tighter bounds across noisy stochastic gradient descent (SGD) steps before expressing the result as an (ϵ, δ) budget. Repeated use of a fixed released DP model is post-processing and does not add to its training privacy loss, because a computation that reads only the release, never the protected data, cannot weaken the guarantee the release already satisfies.

Ten $\epsilon = 1$ releases add up to a basic bound of $\epsilon = 10$.

Privacy accounting depends on three primary composition tools. Simple composition bounds the cumulative privacy loss of k pure-DP mechanisms with bounds ϵ_i by $\sum \epsilon_i$; for approximate DP, their additive slack parameters also accumulate. In suitable parameter regimes, advanced (strong) composition improves the bound using square-root and linear terms while allocating additional slack in the probability inequality. For iterative mechanisms like DP-SGD, the moments accountant and Rényi differential privacy developed next track the loss across steps before converting it to an explicit (ϵ, δ) guarantee, and no accountant is uniformly tighter than every alternative. Organizations must maintain a privacy accounting log for the protected population and halt or redesign further data-dependent releases when the allocated bound would be exceeded.

13.9.2.1 Privacy loss random variable and moments accountant

A more refined approach to privacy analysis tracks the privacy loss random variable, which measures how strongly an observation favors one adjacent dataset over the other through a log-likelihood ratio. For adjacent datasets X, X' and algorithm output o, the discrete-output expression is:

$$\mathcal{L}^{(o)}_{\text{priv},(X,X')} = \ln \frac{\Pr[\mathcal{M}(X) = o]}{\Pr[\mathcal{M}(X') = o]}$$

The privacy loss random variable characterizes the log-likelihood ratio between outputs on adjacent datasets; for continuous outputs, the ratio uses densities rather than point probabilities. A sufficient condition for (ϵ, δ)-DP is that every ordered adjacent pair satisfies $\Pr[\mathcal{L}_{\text{priv}} > \epsilon] \leq \delta$, with the output sampled from the numerator distribution. This tail condition is stronger than the general definition, but it enables composition analysis through moment-generating functions.

The moments accountant technique, introduced for DP-SGD by Abadi et al. (2016), bounds exponential moments of the privacy loss distribution. It retains more information about the mechanism than a single per-step privacy bound while still taking the worst case over adjacent datasets. For mechanism $\mathcal{M}$ with privacy loss $\mathcal{L}_{\text{priv}}$, the moments accountant computes:

$$\alpha_{\mathcal{M}}(\lambda_{\text{mom}}) = \max_{X,X'} \ln \mathbb{E}_{o\sim\mathcal{M}(X)} \left[\left(\frac{\Pr[\mathcal{M}(X) = o]}{\Pr[\mathcal{M}(X') = o]} \right)^{\lambda_{\text{mom}}} \right]$$

for moment order λ_{mom}. Because the total privacy loss across steps with fresh noise draws is the sum of the per-step privacy losses, its exponential moment factors into the product of the per-step moments, and its logarithm adds across steps. If each of k steps satisfies the same log-moment bound, composition gives an upper bound of $k \cdot \alpha_{\mathcal{M}}(\lambda_{\text{mom}})$. Instead of converting every step into a separate tail cutoff and summing those cutoffs, the accountant accumulates the log-moments first and applies the Markov inequality once at the end to yield (ϵ, δ) bounds:

$$\epsilon(\delta) = \min_{\lambda_{\text{mom}}>0} \left[\frac{k \cdot \alpha_{\mathcal{M}}(\lambda_{\text{mom}}) + \ln(1/\delta)}{\lambda_{\text{mom}}} \right]$$

The numerical result depends on the sampling rule, sampling rate, noise multiplier, step count, adjacency relation, and accountant. In DP-SGD, random subsampling at rate q also shrinks $\alpha_{\mathcal{M}}(\lambda_{\text{mom}})$ at each step because a given record falls, on average, into only a q fraction of lots, the randomly sampled groups of examples that DP-SGD processes at each step. For example, Abadi et al. (2016) report that with sampling rate $q = 0.01$, noise multiplier $\sigma = 4$, and $\delta = 10^{-5}$, their moments accountant gives $\epsilon = 1.26$ after 100 epochs and $\epsilon = 2.55$ after 400 epochs, whereas the corresponding advanced-composition bounds are 9.34 and 24.22. These values belong to that experiment and are not reusable constants for another DP-SGD run.

13.9.2.2 Renyi differential privacy and composition

Rényi differential privacy (RDP), introduced by Mironov (2017), provides a convenient representation for composition by tracking Rényi divergence. A mechanism $\mathcal{M}$ satisfies $(\alpha_{\text{R}}, \epsilon_{\text{R}})$-RDP if for all adjacent X, X':

$$\mathcal{D}_{\alpha_{\text{R}}}(\mathcal{M}(X) \| \mathcal{M}(X')) = \frac{1}{\alpha_{\text{R}} - 1} \ln \mathbb{E}_{o\sim\mathcal{M}(X')} \left[\left(\frac{\Pr[\mathcal{M}(X) = o]}{\Pr[\mathcal{M}(X') = o]} \right)^{\alpha_{\text{R}}} \right] \leq \epsilon_{\text{R}}$$

where $\alpha_{\text{R}} > 1$ is the Rényi order. RDP composition is simple. If mechanism $\mathcal{M}_i$ satisfies $(\alpha_{\text{R}}, \epsilon_{\text{R},i})$-RDP, then their composition satisfies $(\alpha_{\text{R}}, \sum_i \epsilon_{\text{R},i})$-RDP. This representation avoids repeatedly applying a generic (ϵ, δ) composition theorem to every training step.

For an unsampled Gaussian mechanism with unit ℓ_2 sensitivity and noise standard deviation σ_{mech}, the RDP guarantee is:

$$\epsilon_{\text{R}}(\alpha_{\text{R}}) = \frac{\alpha_{\text{R}}}{2\sigma^2_{\text{mech}}}$$

Subsampling changes this mechanism. In particular, $q^2\alpha_R/(2\sigma^2_{\text{mech}})$ is only a leading-order approximation in restricted regimes, not a privacy guarantee that can certify a DP-SGD run. A sampled-Gaussian accountant must evaluate the specified sampling scheme and adjacency relation at multiple Rényi orders, compose the resulting per-step bounds across the actual number of steps, and then convert the result to (ϵ, δ)-DP. Production implementations should use a tested accountant such as Google DP Accounting, TensorFlow Privacy, or Opacus and record its version and sampling convention with the released model.

This RDP guarantee supplies a valid (ϵ, δ)-DP bound using the following conversion; the result is an upper bound on privacy loss, not generally its exact value:

$$\epsilon = \epsilon_R + \frac{\ln(1/\delta)}{\alpha_R - 1}$$

optimized over α_R. The moments-accountant and RDP views are closely related; which implementation gives the tightest reported bound depends on the mechanism, orders, and numerical accountant rather than a universal multiplier.

13.9.2.3 DP-SGD accountant inputs for image classification

Consider training a CNN on 50,000 CIFAR-10 images with $(\epsilon, \delta) = (3, 10^{-5})$ target privacy. DP-SGD groups examples into **lots**, the randomly sampled sets whose clipped gradients are summed and share one noise draw per step; with Poisson sampling the lot size varies from step to step around its expected value. The run uses DP-SGD with:

- Poisson sampling with expected lot size $B = 4{,}000$ (inclusion probability $q = B/D = 0.08$)
- Per-example gradient clipping norm $C = 1$
- Noise multiplier $\sigma = 1.3$ (noise scale relative to clipping)
- Training for $N_{\text{epochs}} = 60$ epochs ($k = N_{\text{epochs}} \cdot D/B = 750$ steps)

These inputs do not determine a privacy guarantee through the shortcut $q^2\alpha_R/(2\sigma^2)$. Before release, a tested sampled-Gaussian accountant must evaluate Poisson sampling with add/remove adjacency, compose all 750 steps, report the ϵ it certifies at the specified δ, and verify whether that value meets target $\epsilon = 3$. The release record must pin the accountant library and version, sampling convention, adjacency relation, and evaluated Rényi orders.

The adjacency relation also decides what the bound protects. Per-example clipping bounds the contribution of one record, so when one person contributes many records, as a repeat customer does to a review corpus, a record-level guarantee does not protect that person as a unit. Person-level protection requires bounding each person's total contribution and accounting under person-level adjacency.

Training for $N_{\text{epochs}} = 25$ reduces the number of composed steps, while reducing the expected lot size to $B = 2{,}000$ changes the Poisson inclusion probability to $q = 0.04$. The accountant must be rerun for either change; neither direction should be assigned a numerical ϵ through the leading-order shortcut. If the accountant reports a budget above the target, the operator raises the noise multiplier, trains for fewer epochs, or lowers the sampling rate, and then reruns it. Each change has a utility and compute cost that must be measured. The guarantee itself comes from clipping, calibrated noise, and the privacy accountant; listing 13.1 shows the mechanism sequence without implementing an accountant.

Listing 13.1: DP-SGD Training Step: Per-step clipping, noise addition, and accountant update that together enforce the privacy budget.

```
fix a public normalization B for the expected lot size under the sampling design
check prospectively that the accountant permits the next step within the budget
independently sample each record with probability q (Poisson sampling)
clip each sampled per-example gradient to norm C
sum clipped gradients; use the zero vector for an empty lot
add noise N(0, sigma^2 C^2 I), then divide by fixed B
take one optimizer step with the privatized gradient
record the step in the matching privacy accountant
stop before a step that would exceed the release budget
account for private-data-dependent selection; retuning does not erase prior releases
```

13.9.3 Quantifying the privacy-utility trade-off

Tightening differential privacy to single-digit ϵ keeps MNIST within 3.3 percentage points of its 98.3 percent nonprivate baseline, yet drops CIFAR-10 trained from scratch by roughly 14 to 22 percentage points even at $\epsilon = 8$ (Abadi et al. 2016; Bu et al. 2020; De et al. 2022). That gap shows how strongly the utility cost of differential privacy depends on task complexity, architecture, privacy accountant, pretraining, and whether extra public data are allowed.

Figure 13.26 collects those published MNIST and CIFAR-10 results on a logarithmic axis (Abadi et al. 2016; Bu et al. 2020; De et al. 2022). These comparisons show setting-specific utility costs; differences in architectures and training procedures prevent isolating task complexity alone. The plotted MNIST results stay within 3.3 percentage points of the 98.3 percent nonprivate baseline at every plotted single-digit budget, the smallest of which is $\epsilon = 2$. For CIFAR-10 without extra data, the plotted results reach 73 percent in Abadi et al. (2016) and 81.4 percent in De et al. (2022) at $\epsilon = 8$, roughly 14 to 22 points below the displayed 95.0 percent nonprivate reference. These CIFAR-10 settings therefore pay several times the MNIST accuracy cost at a budget four times looser. That reference is a comparison point, not a ceiling on achievable accuracy. Selecting a budget therefore requires the full training regime, adjacency relation, additive slack, accountant, and utility requirements rather than the horizontal coordinate alone.

These empirical results show why a privacy budget alone does not predict utility. The original DP-SGD experiments achieved higher accuracy on MNIST than on CIFAR-10, while later training methods improved the CIFAR-10 result under the same budget. Public pretraining changes the learning problem again. With an NFNet-F3 pretrained on JFT-4B, De et al. (2022) report 86.7 percent ImageNet accuracy at $(8, 8 \times 10^{-7})$-DP, so comparisons must state which data receive protection and which data support pretraining. The figure supplies selected operating points rather than a complete privacy-accuracy curve. A deployment needs its own sweep over ϵ with the other privacy assumptions and acceptance criteria held explicit.

13.9.4 Decision framework for differential privacy deployment trade-offs

Committing to $\epsilon = 1$ before checking rare-subgroup accuracy or 8-GPU training budgets can leave a clinical classifier too inaccurate or too expensive to ship. The quantified trade-offs turn differential privacy from a mathematical guarantee into a deployment decision: specify the required protection first, then test whether a mechanism can deliver it at acceptable utility and systems cost. A mandatory privacy bound is a feasibility constraint; an accuracy improvement does not by itself justify exceeding it.

Start with the regulatory and legal requirement. DP becomes especially valuable when policy, contracts, audits, or risk assessments require demonstrable privacy bounds. Regulations rarely mandate DP by name, but GDPR's data-protection-by-design principle (European Parliament and Council of the European Union 2016) and privacy-by-design framing (Cavoukian 2012) motivate privacy-preserving system design. HIPAA's technical safeguards, in turn, require controls such as access control, audit controls, integrity protections, authentication, and transmission security (U.S. Department of Health and Human Services 2005). DP can complement those controls by documenting a mathematical privacy-loss bound, but it does not replace HIPAA safeguard implementation. Organizations handling EU health data, US medical records, or California personal

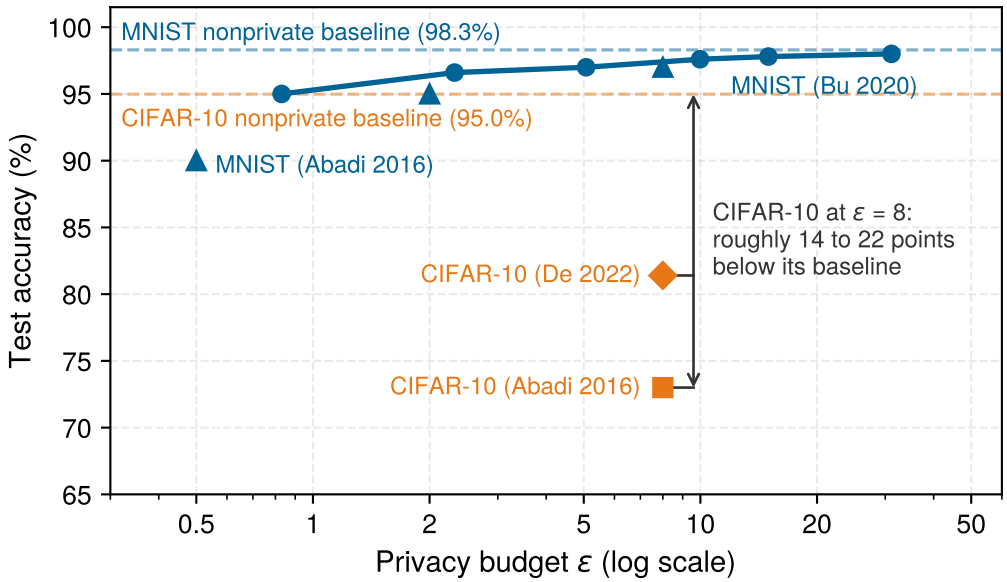

Figure 13.26: Published Privacy-Utility Results: Accuracy at selected privacy budgets from (Abadi et al. 2016; Bu et al. 2020; De et al. 2022). MNIST stays comparatively accurate under single-digit epsilon budgets in these experiments, while CIFAR-10 trained from scratch without extra data remains harder, with De et al. reporting 81.4 percent at $\epsilon = 8$. Comparisons require the full mechanism, accountant, adjacency relation, training regime, and data rather than epsilon alone.

data subject to California state privacy laws (CCPA/CPRA) should evaluate DP as one possible compliance-supporting mechanism.

Next, match the threat model to the mechanism. DP bounds how much any one record can change the released distribution, even when an adversary chooses queries adaptively and combines outputs with auxiliary information. This property can limit membership inference and record-level extraction, but only under the stated adjacency relation and composed budget. When the primary threats are data breaches of stored data rather than inference from a release, encryption and access controls address the relevant boundary; DP does not substitute for them.

The dataset and task then determine whether the guarantee is usable. Effective DP training depends on whether the clipped signal remains learnable after adding calibrated noise. Effective sample size, the number of independent examples that actually inform learning for a class or subgroup, captures why repeated observations, class imbalance, and subgroup rarity can limit learning even in a large table. Task complexity also matters, but class count alone does not determine noise tolerance; representation, label quality, and the training regime affect the result. The data sensitivity distribution raises a separate question about whose errors carry the greatest consequences; sensitivity here means the cost of an error, not the mathematical sensitivity used to calibrate noise. Rare subgroups may suffer disproportionate utility loss even when the same formal privacy bound covers every protected contribution. As the salary notebook showed, the noise that hides one contribution is the same for every group, while a small group supplies little signal to outweigh it.

The utility budget must also be explicit. Before deployment, the team must define task-specific acceptance criteria for overall performance, rare subgroups, calibration, and safety-relevant errors. An average accuracy change that appears small can still be unacceptable if DP noise concentrates errors on a protected subgroup or a high-cost clinical class. Published results supply comparison points on the privacy-utility frontier (figure 13.26), but the release decision requires an experiment on the actual model and data.

Finally, the computational budget decides whether the design can ship. DP-SGD adds clipping, noise generation, and accounting work, while attaining acceptable private utility can change the batch size and training duration. The resulting runtime and memory cost depends on the model

and implementation. Deployment planning needs measured time to the required utility under the privacy budget, rather than a universal compute multiplier.

13.9.4.1 Decision matrix summary

Use DP when a formal participation or contribution bound addresses the stated threat model, a tested accountant covers the relevant data-dependent releases, measured utility satisfies the task's acceptance criteria, and the training system can afford the mechanism's compute and memory cost. Use different controls when the primary boundary is unauthorized access to stored data. If a required DP bound cannot coexist with acceptable utility, redesign the mechanism, task, or data strategy, or withhold the release. Federated learning, encryption, access control, and audit logging do not substitute for a required release guarantee merely because they are easier to deploy.

Checkpoint 13.4: Applying the DP decision framework

A wearable cardiac monitor classifies arrhythmias from raw ECG signals collected on-device for 5,000 enrolled patients. The product is sold in the United States and is HIPAA-regulated; the training pipeline runs on a cloud cluster of 8 GPUs, and a competing vendor has incentive to extract the model. Apply each of the five decision criteria to determine whether DP-SGD belongs in the training pipeline.

- ☐ **Criterion 1 (regulatory)**: decide whether HIPAA supports a DP guarantee here and identify the documentation DP supplies that ad-hoc de-identification cannot.
- ☐ **Criterion 2 (threat model)**: identify the adversary populations that matter for a cardiac-monitor model and evaluate membership inference as a credible attack.
- ☐ **Criterion 3 (dataset characteristics)**: evaluate effective sample size, class imbalance, repeated measurements, and rare clinical subgroups before selecting a mechanism.
- ☐ **Criterion 4 (acceptable utility loss)**: define safety-critical error, calibration, and subgroup acceptance criteria, then measure whether the private model satisfies them.
- ☐ **Criterion 5 (compute)**: estimate the runtime and memory cost of private training on the stated cloud allocation, measure utility under the selected privacy budget, and identify the binding deployment constraint.

This framework treats DP as one tool in a privacy-preserving toolkit rather than a universal solution. Production deployments often combine an accounted DP mechanism with data minimization, federated execution where raw-data centralization is out of scope, secure aggregation for contribution confidentiality, and access controls for deployment security.

Returning to the data-layer techniques of section 13.8.5, differential privacy sits beside federated learning, homomorphic encryption, secure multiparty computation, and synthetic data generation, and each strikes its own balance among privacy strength, accuracy, and computational cost (table 13.8). The maturity labels are a deployment snapshot for the regimes discussed here, not a permanent ranking of the underlying techniques. Cryptographic guarantees depend on key custody and the stated adversary model; they do not prevent an authorized output from revealing sensitive information. Differential privacy instead bounds disclosure through the released result under a specified adjacency relation.

Table 13.8: Privacy-Preserving Technique Comparison: Data privacy techniques impose varying computational costs and offer different levels of formal privacy guarantees, requiring practitioners to balance privacy strength with model utility and deployment constraints. Weighing privacy guarantees, computational overhead, maturity, typical use cases, and trade-offs together guides the design of privacy-aware machine learning systems.

Technique	Privacy guarantee	Computational overhead	Deployment snapshot	Typical use case	Trade-offs
Differential Privacy	Formal (ϵ, δ)-DP	Moderate to High	Production	Training with sensitive or regulated data	Reduced accuracy; careful tuning of ϵ/noise required to balance utility and protection
Federated Learning	Local data retention only	Moderate	Production	Cross-device or cross-org collaborative learning	Gradient leakage risk; requires secure aggregation and orchestration infrastructure
Homomorphic Encryption	Ciphertext confidentiality	High	Experimental	Inference in untrusted cloud environments	High latency and memory usage; suitable for limited-scope inference on fixed-function models
Secure MPC	Protocol-dependent input privacy	Very High	Experimental	Joint training across mutually untrusted parties	Expensive communication; challenging to scale to many participants or deep models
Synthetic Data	None from synthesis alone	Low to Moderate	Emerging	Data sharing, benchmarking without direct access to raw data	May leak sensitive patterns if training process is not differentially private or audited for fidelity

Increasing the noise to reduce ϵ may degrade model accuracy, especially in low-data regimes or fine-grained classification tasks. Consequently, DP is often applied selectively (either during training on sensitive datasets or at inference when returning aggregate statistics) to balance privacy with performance goals (Dwork and Roth 2014).

Differential privacy completes the data layer of the defense stack. With the mechanisms of every layer now in view, the remaining engineering questions are which of them a given deployment needs and in what order to adopt them.

13.10 Selecting and Sequencing Defenses

Moving from mathematical guarantees to production controls requires more than setting an epsilon value. A common mistake is attempting to deploy every defense at once, including differential privacy, trusted execution environments, adversarial training, model watermarking, and red-team automation, while basic controls such as access boundaries and encrypted transport remain weak. Two decisions prevent that mistake. The first selects the defenses a deployment context requires; the second sequences them so that each control rests on the ones beneath it.

13.10.1 Defense selection by deployment context

A hospital training a federated diagnostic model on patient records and a public-facing customer chatbot face completely different attackers, latency budgets, and legal mandates. The hardware mechanisms of table 13.7, the deployment and runtime controls, and the cryptographic and statistical privacy techniques of the data layer therefore form a menu rather than a fixed recipe. The appropriate defense for any given deployment depends on three interacting factors: the threat model (who attacks and with what capabilities), the deployment context (what computational and latency budgets exist), and the regulatory environment (what legal mandates constrain design). A defense selection framework maps seven common deployment contexts to their primary threats, recommended defense combinations, and the trade-offs each combination requires evaluating (table 13.9). The rows that include differential privacy require an accounted mechanism, built from the sensitivity and composition machinery of section 13.9, rather than a universal privacy-budget value.

Table 13.9: Defense Selection Framework: Maps deployment contexts to threat-specific defensive strategies and deployment trade-offs. The framework provides starting points for security architecture design, highlighting primary threats, recommended defense combinations, and key implementation trade-offs across seven common ML system deployment scenarios.

Deployment context	Primary threats	Recommended defenses	Key trade-offs
Healthcare ML (federated diagnostic models)	Data leakage (HIPAA violation); membership inference; unauthorized access	Accounted differential privacy for training when the threat model requires it; federated learning across hospitals; TEE for inference on sensitive data; audit logging and RBAC	Privacy and TEE parameters require task-specific utility and latency measurement
Financial ML (fraud detection API)	Model theft (IP loss); adversarial evasion; data poisoning	Model encryption (AES-256) at rest; HSM for cryptographic key management; adversarial training on perturbed transactions; input validation and per-principal rate limiting; output confidence monitoring	HSM units cost roughly $20,000–$100,000 each before integration and certification; rate limiting may impact legitimate high-frequency users
Edge ML (mobile/IoT devices)	Physical access; side-channel attacks; model extraction	Secure boot; a suitable TEE; authenticated encrypted model storage; PUF-backed device identity	Protected working set, paging support, power cost, and physical-attack resistance depend on the platform
Cloud ML Training (multi-tenant platform)	Data poisoning; backdoor injection; gradient leakage	Secure data pipelines with provenance; accounted DP-SGD when participation leakage is in scope; gradient verification and anomaly detection; secure aggregation for federated training; model watermarking for IP protection	DP and verification overhead depend on implementation and workload; federated aggregation requires authenticated, confidential communication
Public-Facing LLM (chatbot/API)	Prompt injection; data extraction (training leakage); abuse/overuse	Instruction-data separation; input/output policy checks; isolated execution; authentication and rate budgets	Measure latency and false-positive cost; filtering and confidence scores do not guarantee resistance to injection
Multi-Party ML (cross-organizational training)	Data sharing restrictions; curious or malicious participants; regulatory duties	Federated learning (no raw data sharing); SMPC for secure aggregation; accounted differential privacy when participation leakage is in scope; homomorphic encryption for selected operations	Communication and cryptographic overhead depend on topology and protocol; utility must be measured for the selected privacy mechanism; requires agreements for liability
Critical Infrastructure (autonomous vehicles, power grids)	Supply chain compromise; real-time adversarial attacks; safety-critical failures	Hardware attestation (trusted platform module, or TPM, and PUF); Secure Boot and runtime integrity checks; redundant model validation; fault injection detection; fail-safe fallback mechanisms	Validate physical threat coverage, latency, fail-safe behavior, and the required assurance process

Checkpoint 13.5: Mapping hardware and deployment threats to defenses

The chapter introduced six hardware-level attack categories, four hardware-security primitives (table 13.7), and a framework mapping seven deployment contexts to full-stack defenses (table 13.9). Apply that combined mapping to the scenarios below.

- ☐ **Primitive matching**: match fault injection on model weights and side-channel extraction from the inference pipeline to the most relevant hardware-security primitives and runtime controls across the critical-infrastructure and edge rows of table 13.9.
- ☐ **Edge defense stack**: select a combination of hardware and deployment controls for an edge inference deployment facing physical access, side-channel attacks, and model extraction.
- ☐ **Supply-chain gap**: identify which supply-chain compromises remain outside the four hardware primitives' guarantees in the selected edge or cloud-training deployment, and specify the provenance or inspection evidence needed to assess those risks.
- ☐ **TEE vs. attestation**: distinguish runtime isolation from cryptographic evidence about an execution environment, and identify which fault-injection, counterfeit-hardware, and output-leakage risks remain outside those guarantees.

A recurring theme across table 13.9 is the cost of privacy. Differential privacy appears in three of seven deployment contexts, and each such row exposes an accuracy or systems cost that must be weighed against regulatory mandates and risk tolerance. The table is therefore a deployment triage tool rather than a universal recipe, because serialization, artifact integrity, access control, runtime isolation, and privacy budgets each defend a different boundary. Selecting the right defenses for a context settles what to build; the maturity model settles the order in which to build it.

13.10.2 Security and privacy maturity model

Deploying homomorphic encryption or a red-team harness while model checkpoints sit in an unauthenticated bucket and API keys lack per-principal budgets leaves the easiest attack paths wide open. A maturity model asks which system property is currently unprotected: access boundary, data privacy, model integrity, adversarial robustness, or governance evidence. The dependency structure in table 13.10 keeps that choice tied to the boundary being defended.

Table 13.10: ML Security and Privacy Maturity Model: Security controls should be added according to the threat model and deployment context. The order is not a calendar; it is a dependency structure from basic access boundaries to stronger privacy, integrity, adversarial, and governance controls.

Maturity layer	Primary question	Control family
Access and configuration boundary	Who can touch data, weights, deployments, and logs?	Least privilege, service identity, encrypted transport, audit logging
Data and privacy boundary	What can be learned from training, fine-tuning, or output?	Privacy accounting, output limiting, secure aggregation, sensitive-data isolation
Model integrity boundary	How is the validated model and data path authenticated?	Signed artifacts, registry controls, hash checks, release gates, rollback paths
Adversarial and abuse boundary	What can an adaptive attacker extract, infer, or induce?	Query monitoring, rate limiting, robustness evaluation, red-team exercises
Governance and evidence boundary	Which obligations must be proven after deployment?	Compliance mapping, incident records, retention policy, reproducible control checks

The maturity layers are cumulative. Access controls and encrypted channels close paths that model-interface defenses do not cover, such as an exposed registry. Privacy analysis begins with the original training data and continues through fine-tuning, personalization, and telemetry feedback loops. When DP is required, the release budget becomes a system resource covering the relevant training, selection, and evaluation mechanisms that access protected data. Model-integrity controls protect the artifact path: signed manifests, key management, hash checks, and release gates connect the loaded model to the approved version, subject to the enforcement boundary remaining intact.

Adversarial evaluation tests deliberate, adaptive behavior in addition to accidental failures. Query monitoring and output limiting constrain extraction attempts, while robustness evaluation and red-team exercises test whether an attacker can induce unsafe behavior. Trusted execution environments and certified robustness defenses protect different boundaries, so each needs a threat model that justifies its costs and assumptions. Governance defines obligations and evidence requirements from the start, rather than entering only after technical controls. Audit logs, compliance mappings, incident records, and reproducible checks then help assess whether those obligations remain satisfied as the model, data, and traffic change.

Different domains traverse the model differently. A healthcare deployment prioritizes privacy accounting and governance evidence before broad patient-facing use. A financial fraud system emphasizes access control, model integrity, and abuse monitoring because extraction and evasion directly change loss exposure. An autonomous system emphasizes adversarial robustness and incident evidence because safety depends on behavior under distribution shift and physical-world perturbations. The point is not to follow a universal implementation calendar; it is to make the threat model determine which control boundary must mature next.

The evidence produced by these controls supports a separate governance decision: which uses are permitted, who accepts residual risk, and who can require remediation. Section 16.2 develops those institutional responsibilities. Here the engineering obligation is to make access decisions, release

budgets, and incident histories verifiable. By sequencing controls around those boundaries, the maturity model moves security from reactive patching toward proactive threat control, though even well-designed systems can still fail when teams mistake a local mechanism for a complete guarantee.

13.11 Fallacies and Pitfalls

A common ML security failure is treating a local mechanism as a system guarantee. Obscurity, differential privacy libraries, federated learning, or encrypted storage each solve one part of the threat model, but none substitutes for lifecycle-wide reasoning about data, models, interfaces, hardware, and distributed scale.

Fallacy: *Security through obscurity provides adequate protection for machine learning models.*

Hiding architectures or parameters does not establish security when an attacker can learn functionality through an exposed API. The extraction examples in section 13.5.1.4 show why query budgets, returned information, and the target task must be evaluated together; no universal query count or accuracy threshold applies. Adversarial examples can also transfer from surrogate models, although success depends on the models, data, and perturbation constraints (Nicolas Papernot, McDaniel, and Goodfellow 2016). Kerckhoffs's principle motivates defenses that remain useful when the mechanism is known. Access controls and query limits constrain exposure, output restrictions reduce available information, and monitoring supports investigation; watermarking serves attribution rather than preventing copying.

Pitfall: *Assuming that differential privacy automatically ensures privacy without considering implementation details.*

Treating differential privacy as a universal solution without specifying adjacency, parameter selection, and budget tracking leaves the release claim incomplete. The same ϵ can describe different protections under different adjacency relations, values of δ, sampling schemes, and compositions, while utility depends on the model, task, and data. Privacy budgets compose across additional data-dependent releases: releasing 10 separately trained models at $\epsilon = 1.0$ each gives a basic-composition bound of $\epsilon = 10$, not $\epsilon = 1.0$, although a mechanism-specific accountant may provide a tighter bound. Retraining, tuning, and evaluation mechanisms that access protected data must be included in the accounting boundary.

Fallacy: *Federated learning inherently provides privacy protection without additional safeguards.*

This misconception assumes decentralization automatically ensures privacy, but individual updates can expose membership or reconstruct features of local training data. As examined in section 13.8.5.1, effective federated privacy requires controls matched to the adversary: secure aggregation hides individual contributions from the coordinator, an accounted DP mechanism bounds record- or user-level participation leakage under its stated adjacency relation, and authenticated encryption protects updates in transit. Organizations cannot claim privacy merely because raw records remain on devices; the protocol still releases information derived from those records.

Pitfall: *Planning poisoning defenses only for large training-data compromise.*

A small, strategically chosen set of poisoned examples can target a decision without causing a conspicuous drop in overall validation accuracy (Shafahi et al. 2018). Backdoor attacks make this distinction explicit. A model may behave normally on clean inputs while producing attacker-chosen behavior when a trigger appears (Gu et al. 2017). The required poisoning fraction and resulting attack success depend on attacker access, the training procedure, and the trigger; a universal percentage would conceal those dependencies. Defenses must therefore evaluate targeted behavior and provenance, including user-generated ingestion paths, rather than treating a low estimated contamination rate as a safety guarantee; section 14.7 develops them. Retraining or transfer learning also requires renewed testing instead of an assumption that the trigger has disappeared.

Fallacy: *Security can be isolated from the rest of the ML system.*

Component defenses leave gaps when the same asset remains exposed through another path. Per-user query quotas, output perturbation, and top-k filtering at a public API do not protect weights in an unauthenticated registry or query-response pairs in exposed logs. As established in section 13.4.1,

threat modeling must follow the asset through collection, training, storage, deployment, and monitoring. A failed boundary can defeat controls that depend on it, while independent controls may still limit the compromise. The engineering task is to identify both the bypass path and the protections that remain effective, rather than assume either that one protected interface secures the whole system or that every local failure defeats every guarantee.

Pitfall: *Underestimating the attack surface expansion in distributed ML systems.*

Organizations secure single-node ML systems effectively but fail to recognize that distributed architectures introduce additional trust boundaries and communication channels. Each link is a channel whose endpoints must authenticate each other and protect the traffic between them, even when they reuse node-level credentials. An all-to-all topology has $\mathcal{O}(N_{\text{nodes}}^2)$ links, whereas a ring has $\mathcal{O}(N_{\text{nodes}})$ links. Link count alone does not determine vulnerability count, but every additional node, credential, and channel must be included in the threat model.

In the reference training cluster of 1,024 GPUs (128 nodes with 8 GPUs each), every node is at once a gradient contributor, a fabric endpoint, and a key holder. Section 13.6 therefore treats worker admission, collective traffic, shared accelerators, and attestation as boundaries of their own. Federated learning with 10,000 clients creates 10,000 separately managed endpoints whose attestation support and coverage may differ. Aggregation methods that inspect individual updates also need reconciliation with protocols that hide those updates; combining the two requires a design that preserves both guarantees. Effective distributed ML security therefore requires threat modeling the added nodes and boundaries: securing inter-node channels, managing identity across security domains, and coordinating policies across heterogeneous infrastructure.

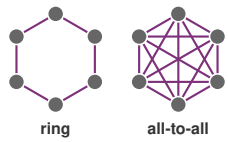

Same nodes, many more channels, and every all-to-all link joins the threat model.

Recognizing that the shift from a single node to a distributed cluster adds nodes, channels, credentials, and trust boundaries rules out perimeter-only thinking. Security for the machine learning fleet must be layered, measured, and tied to the full lifecycle.

13.12 Summary

Privacy and security protect the machine learning fleet's models, data, hardware, and service boundaries from extraction, poisoning, manipulation, and unauthorized access. The breach cases establish the boundaries a fleet must defend: Strava shows that an aggregate release can disclose behavior without leaking a single record, Stuxnet exposes supply-chain trust, the Jeep Cherokee hack exposes insufficient isolation, and Mirai exposes weak endpoint identity. Learned behavior adds further boundaries because poisoned data, extraction queries, and prompt injection can compromise a system through otherwise legitimate interfaces. The fleet adds boundaries of its own that no API gateway sees: the gradient each worker contributes, the fabric that carries it, the accelerators tenants share, and the evidence that each node runs approved code. Hardware trust anchors protect execution, provenance connects artifacts to release decisions, runtime controls contain abuse, and differential privacy bounds participation leakage. These mechanisms compose only when each is tied to the actor and exposure it addresses.

Their costs span all three C^3 axes and can enter different terms of the Fleet Law (equation 1.3). Output shaping, input validation, redundant gradient checks, DP-SGD clipping and noise, and protected execution can add work to T_{compute}. Fabric encryption, gathering individual updates for some robust aggregation methods, and masked federated uploads can increase $T_{\text{comm}}(N)$; offload and overlap determine how much reaches the critical path. Admission checks, key rotation, privacy-budget checks, and aggregation thresholds can add to $T_{\text{sync}}(N)$ when the workflow waits for them. Federated protocols also need deadlines and dropout handling rather than an assumption that every client responds. Capacity planning must measure these costs in the deployed path so that operating targets do not depend on bypassing a control under load.

Key Takeaways: Defend the model, not just the server

- **The model is the attack surface**: ML systems inherit ordinary server threats but add attacks against the learned boundary itself. Poisoned data, model extraction, adversarial

inputs, and prompt injection can compromise behavior without exploiting a traditional software vulnerability.

- **Extraction is bounded by cumulative budgets**: A daily rate limit slows a query-based extractor without capping what it collects. A lifetime per-principal budget priced at break-even with retraining, paired with an output contract that limits what each response reveals, removes the cost advantage of copying, but only as far as principals map to real actors.
- **Privacy is an accounted budget**: Because models can encode information about sensitive data, a correctly executing model can disclose training data through memorization and membership inference, so privacy cannot be treated as anonymization at the perimeter. Differential privacy (principle 23) offers bounded (ϵ, δ) claims only when all data-dependent DP mechanisms are composed correctly; querying a fixed DP-trained model is post-processing and does not add training privacy loss.
- **Generative systems need enforceable control boundaries**: Prompt injection, jailbreaks, and PII leakage pass through normal text channels. Treat model output as untrusted, keep instructions separate from untrusted content, and keep authorization in an external policy layer.
- **The fleet's own interfaces are trust boundaries**: An unbounded malicious update can dominate a mean; the examined gather-based implementation of robust aggregation adds about $N_{\text{DP}}/2$ times the traffic, an implementation cost rather than a universal lower bound. Worker admission, collective traffic, and shared accelerators need separate controls, and admission does not replace update checks, because an admitted worker can still submit a malicious update.
- **Provenance bounds trust and enables recovery**: A model can be trusted only as far as its record of data, code, weights, and hardware reaches (principle 22). Signed datasets, locked dependencies, registry permissions, and deployment attestations let operators answer which data, code, principal, node, and checkpoint produced a suspect model. Without that chain, rollback and forensics become guesswork at fleet speed.
- **Hardware trust has explicit boundaries**: TEEs isolate execution, secure boot verifies an authorized boot chain, and HSMs protect keys and cryptographic operations. These mechanisms address different trust boundaries; they do not establish that model behavior is correct or that an authorized output preserves privacy.

A fleet can execute every authorized operation correctly and still disclose sensitive information through an approved release. The information leakage invariant (principle 23) therefore makes the release boundary as important as the execution boundary. A deployment decision must establish both who can access the computation and what its outputs allow an observer to learn.

What's Next: From security to resilience

Security and privacy controls address adversarial intent and control, establishing who can infer, manipulate, disable, or authorize what across data, model, interface, and infrastructure boundaries. Those controls decide who may reach the model and what its releases may disclose. They cannot show that its answers stay valid once inputs, data distributions, or hardware depart from the conditions under which it was validated. A fleet can pass every access check and every health signal while its predictions degrade without raising an error, whether from road conditions its training data never covered, from a crafted input that satisfies the interface contract, or from a faulty accelerator that keeps returning plausible numbers. Detecting and bounding that silent failure requires behavioral evidence about accuracy, calibration, and fallback under environmental shift, input-level attack, and system faults, which is the robustness problem the next chapter takes up.

↪ Drift detection, adversarial-defense evaluation, and poisoning defenses are developed in chapter 14.

14

Robust AI

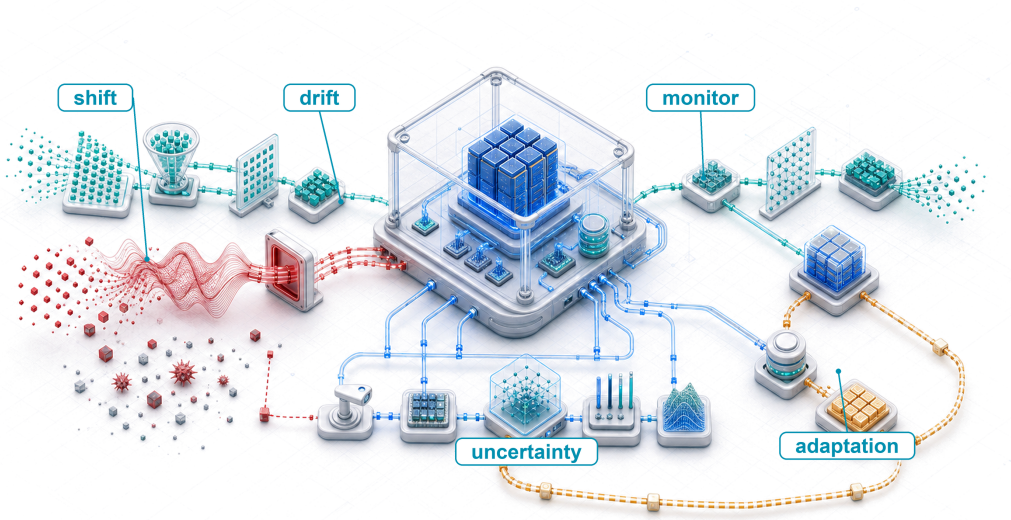

14.1 The Silent Failure Problem
14.2 Real-World Robustness Failures
14.3 A Unified Framework for Robust AI
14.4 Environmental Shifts
14.5 Adversarial Attacks
14.6 Adversarial Defenses
14.7 Data Poisoning Defenses
14.8 Robustness of Generative Models
14.9 System-Level Faults
14.10 Fallacies and Pitfalls
14.11 Summary

Purpose

Why can machine learning systems fail silently even when conventional health signals remain normal?

A model service can maintain stable latency and uptime while its predictions become unreliable. Changes in operating conditions, manipulated inputs, or corrupted computation may leave ordinary health signals untouched, allowing errors to affect decisions long before business metrics expose them. Conventional software can also fail silently; machine learning adds statistical failures that a successful request cannot reveal. At fleet scale, replicating a model across healthy machines does not remove a weakness shared by every copy. Traffic also spans different regions, devices, and user populations, so acceptable average performance can conceal a serious failure in one group. The decision to continue answering therefore requires evidence about the conditions under which the predictions remain useful, not merely confirmation that the service is available. Producing that evidence consumes resources, and responding to it creates further obligations. Additional checks can delay answers; rejecting uncertain requests can overwhelm reviewers or a backup service. A replacement model must demonstrate acceptable behavior before the fleet adopts it, while waiting too long can prolong exposure to errors. The engineering challenge is to connect evidence to a bounded response without exhausting the resources needed to serve useful predictions. Protection must cover the model and the surrounding data path, because either can invalidate the result. In C^3 terms, verification spends compute and communication, while coordination determines when the evidence should trigger investigation, fallback, or a halt.

Learning Objectives

- Classify a silent failure by robustness pillar and diagnose its source using input replay, deployment history, and system telemetry
- Distinguish covariate shift, label shift, and concept drift, and judge when drift evidence justifies a model update
- Apply a risk-coverage curve to choose a selective-prediction threshold and size its fallback load, and contrast it with calibration and with conformal coverage under exchangeability
- Evaluate adversarial attacks by attacker access, and select defenses such as adversarial training, certification, and input detection
- Select layered poisoning defenses that combine source controls, anomaly detection, robust training, and post-training checks
- Evaluate generative-model reliability against a named evidence source, treating entropy, self-consistency, and output guardrails as priced routing signals
- Evaluate robustness designs against clean accuracy, compute, latency, and fallback-capacity budgets

14.1 The Silent Failure Problem

A vehicle's vision model validated on sunny roads may become unreliable when snow changes the visual distribution, yet reading a snow-covered stop sign as a speed-limit sign or misclassifying a truck as the sky need not throw an exception or produce a stack trace. A demand forecasting model does not error out when it produces wildly inaccurate predictions, and a medical diagnosis system does not shut down when it quietly returns incorrect classifications that endanger patient lives. Because these errors appear in predictions while infrastructure telemetry remains green, this **silent failure mode** requires ML-specific behavioral checks alongside conventional software checks. Whether a prediction error leads to an unsafe action depends on the surrounding perception, planning, and fallback systems, so robustness engineering evaluates and constrains behavior under specified operational stresses, including distribution shift, sensor noise, adversarially crafted inputs, and hardware or software faults that leave the service apparently healthy while its answers become unreliable.

Robust AI therefore occupies the Assurance row of the fleet stack (figure 1.4), where the fleet establishes that its results can be trusted. That row sits directly above Operations at the base of the governance band, and robustness shares it with security and privacy. Security and privacy establish who can manipulate, extract from, or infer through the model and which controls constrain that access or disclosure. Robustness adds a behavioral requirement when inputs, data distributions, hardware, or software depart from the conditions assumed during validation, because only a behavioral test reveals how a model responds outside its training distribution. Within a specified operating envelope, the fleet needs evidence that predictions or fallback actions remain acceptable. Adversarial examples, inputs crafted to force an error, become perturbations to measure and defend against, with certification available under particular assumptions. Nonadversarial drift and faults require their own evaluations. In operation, diagnosis of the failure source comes first, using the signals developed in section 14.9.2. Population-shift tests then supply evidence of environmental change, while attack-specific evaluation tests whether deliberate manipulation defeats the model or its defenses.

Checkpointing and recovery restore execution state, while access control enforces permissions at the system boundary. These mechanisms do not establish that predictions remain useful. An adversarial input may satisfy the interface contract, the data distribution may drift beyond the model's validated conditions, or a preprocessing fault may alter inputs without stopping inference. Recovery can restore a clean checkpoint after some failures, but behavioral checks must determine whether that restored system meets its acceptance criteria. Fault detection, diagnosis, and recovery therefore need evidence about both execution and prediction quality, motivating a precise definition of Robust AI:

Definition 14.1: Robust AI

Robust AI is the measurable systems property that a model's predictions remain valid (within specified error bounds) under distribution shift, adversarial perturbation, and hardware or software faults, as opposed to the average-case accuracy achieved under ideal independent and identically distributed (i.i.d.) conditions.

1. **Significance**: Robustness is quantified against an explicit operating or threat set. A certified classifier can prove that its prediction cannot change for any input within a specified perturbation set, such as an ℓ_∞ ball of fixed radius around one test point (every input coordinate changed by at most that radius). Empirical stress tests instead measure performance on named corruptions, shifts, or attacks without proving behavior outside the evaluated cases.
2. **Distinction**: Standard generalization measures average-case accuracy on held-out i.i.d. data drawn from the training distribution. Robustness measures behavior under specified departures from that distribution, including natural shifts, semantic transformations, physical changes, malicious queries, and norm-bounded perturbations. Some evaluations are worst-case within a formal set; others provide empirical evidence over a sampled deployment envelope.
3. **Common pitfall**: A frequent misconception is that robustness can be added as a post-hoc monitoring layer to any existing model. Monitoring supplies evidence about observed behavior; it does not establish certified robustness. Training-time hardening and serving-time defenses each require evaluation of the complete defended system against a specified threat model. A certificate applies only within its stated perturbation set and assumptions.

Two models can agree on every clean input and still differ in how far an input can move before its label flips. In figure 14.1, each labeled point is surrounded by the same perturbation radius, the largest input change the evaluation allows. Vertical boundary A keeps every allowed change on the correct side, while oblique boundary B cuts through the neighborhood around x_+, so a small displacement δ_{adv} flips the predicted class. The adversarial defenses of section 14.6 turn that distance into a measured quantity and, under stated assumptions, a certified one.

Three categories organize these silent failures and the responses they require, and the rest of the chapter calls them the **three pillars** of robust AI. The first, environmental change, includes distribution shifts, concept drift, and evolving operating conditions. A model trained on last year's transaction patterns may become unreliable as customer behavior evolves. Monitoring must establish whether the change affects predictive quality before an adaptation strategy can address it.

The second category, input-level attacks, is malicious manipulation that targets model behavior directly through the data a model receives or learns from. Adversarial attacks, data poisoning attempts, and prompt injection seek incorrect predictions or unauthorized behavior. Authentication and access controls constrain who can submit data and which actions the system permits, but an authorized input can still induce an incorrect result. Behavioral evaluation must therefore examine the manipulation possible within the access that remains.

The third category is system-level faults: hardware faults, software bugs, dependency failures, and runtime errors that corrupt the machinery around the model. Software and pipeline defects belong to this pillar by origin, yet their symptoms can amplify, mask, or mimic the other two, which is why they also cut across the whole taxonomy. Bugs, design flaws, and implementation errors within algorithms, libraries, and frameworks propagate through the system, creating systemic vulnerabilities[1] that transcend individual component failures. A preprocessing bug might create artificial distribution shifts; a numerical error might corrupt model behavior in ways indistinguishable from adversarial attack; a race condition might corrupt learned representations. Because these faults originate in the systems layer, their detailed taxonomy and mitigation strategies are covered in chapter 7; this chapter focuses on how system-level faults interact with environmental shifts and input-level attacks.

[1] **Systemic vulnerability**: A shared dependency can expose many components to the same defect or attack. Pinning one CUDA or PyTorch version across a fleet makes that version a common dependency, but exploitation still depends on the affected code paths, configuration, and attacker access. Dependency inventories and staged updates help identify and contain the exposed population; replication alone does not remove a shared weakness.

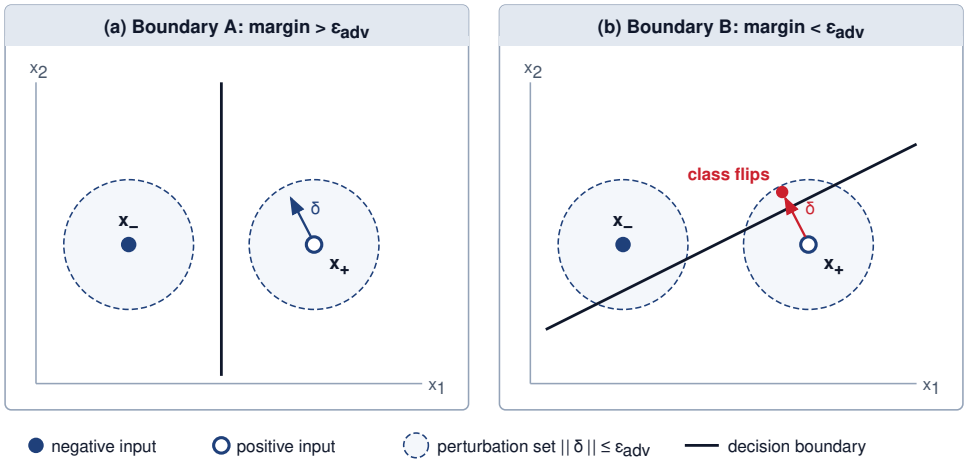

Figure 14.1: Clean Predictions and Perturbation Robustness: Two illustrative two-dimensional classifiers demonstrate the distinction between clean-data accuracy and certified robustness. Both decision boundaries correctly separate the labeled clean inputs. Dashed circles delimit identical perturbation neighborhoods around each point. Boundary A maintains separation outside both circles, whereas boundary B intersects them, allowing a small adversarial perturbation to change the classification. Agreement on clean inputs does not establish resilience under attack.

The appropriate defense depends on where the system runs and which failures it must contain. A cloud service may have capacity for redundant execution that exceeds an edge device's power and memory budget, but additional replicas do not remove a model weakness shared by every copy. Edge devices can concentrate checks on critical inference paths and use a validated local fallback. Both settings require measured protection within their resource and response-time limits.

Each protection consumes resources in the path where it operates. Error correction adds storage and checking work, redundant processing repeats computation, and monitoring uses compute and telemetry capacity. Their effects on bandwidth, energy, and response time depend on the implementation and workload. The robustness decision compares those measured costs with the failures detected or contained and the consequences of errors that remain. Because every layer of the pipeline, from input validation and training through drift detection and fault isolation, carries part of this cost, robustness constrains the architecture rather than wrapping a finished system, and ignoring that constraint has produced severe failures across cloud, edge, and embedded deployments.

14.2 Real-World Robustness Failures

A shared storage outage stops a service outright, whereas a perception or classification error leaves it running with an unsafe result. Separating those two outcomes ensures that a successful request or an operational processor is never mistaken for evidence of a valid prediction. While deliberate input-level attacks are best documented through the controlled field studies of section 14.5, operational failures in the field expose the other pillars across three deployment tiers. Cloud fleets reveal silent hardware corruption during training and harmful errors on population slices, subsets of inputs such as a demographic group or an operating condition, hidden inside aggregate accuracy. Edge vehicles show perception failing outside its validated operating envelope with no time for a remote fallback. Embedded controllers show the hazard a learned component inherits when a false sensor signal drives an irreversible physical action. Even in a healthy cloud service, a model can fail silently on a critical population slice while its aggregate benchmark metrics look strong.

War Story 14.1: The label that exposed the test gap (2015)

Context: In June 2015, Google Photos used automatic image tagging to organize user photos (Kasperkevic 2015).

Mechanism: A deep convolutional image classifier placed Jacky Alciné's photos of himself and a Black friend under the album tag "Gorillas," a severe slice-level failure mode. The failure demonstrates why aggregate top-1/top-5 accuracy over a full test set cannot establish acceptable behavior for high-risk population slices.

Impact: The high-visibility misclassification caused public outrage and exposed a major gap in multi-class vision model safety validation.

Fix: Google disabled the "gorilla" label, blocking the symptom without demonstrating that the classifier itself had been corrected.

Systems lesson: Production ML robustness cannot be defined by aggregate test set accuracy. High-risk classification tasks require slice-level evaluation and a serving response to harmful errors. Automated toxicity guards and validated fallbacks are candidate controls whose coverage must be tested against the specific failure.

14.2.1 Cloud infrastructure failures

Cloud services rest on a hidden assumption that a shared dependency remains both available and correct. The 2017 AWS S3 outage shows the familiar half of that assumption, in which a shared storage dependency failed and otherwise healthy services stopped.[2] (Amazon Web Services 2017) An outage at least announces itself. The robustness problem is the other half, when the dependency stays up but returns plausible, corrupted data, so the service keeps answering while incorrect results propagate.

Silent data corruption (SDC), in which hardware returns a wrong value without raising any error, is that case. The decompression incident analyzed in section 7.5.1 demonstrates the failure signature. A defective CPU returned zero during a size computation for a valid file, causing the decompression pipeline to omit files from its output even though system and kernel logs showed no fault (Dixit et al. 2021). The same section's training illustration (figure 7.7) traces how unhandled corruption can enter gradients, weights, and later checkpoints. In distributed ML training, a corrupted gradient or activation can perturb optimizer state, spread to every replica through gradient synchronization, and persist into later checkpoints. Its blast radius therefore grows with the synchronization and checkpointing design, and recovery may require rollback to the last clean checkpoint (Dixit et al. 2021; Ma et al. 2025). In both settings, checking that execution continues is insufficient; evaluation must test the validity of the resulting data or model behavior.

[2] **AWS S3 outage (2017)**: A mistyped command during routine maintenance removed too much capacity from S3's index and placement subsystems in US-East-1. While those subsystems restarted, S3 could not service requests and dependent AWS services experienced elevated errors or impaired functionality. The incident exposed a single-region dependency pattern: systems that assume regional storage availability as an invariant can fail even when their own application code and model-serving logic remain unchanged.

↩ Silent data corruption during training and its fleet-scale evidence are established in section 7.5.2.

14.2.2 Edge and embedded failures

A cloud training job that detects silent corruption can roll back to a clean checkpoint, whereas an edge system acting in the physical world cannot undo a control decision once issued. In May 2016, a Tesla Model S operating on Autopilot failed to identify a tractor-trailer crossing its path and struck it at highway speed with neither warning nor braking (National Transportation Safety Board 2017). The system had not been validated for that scene, and nothing in its own behavior signaled that it had left its validated envelope. Edge deployments must respond within local compute, power, and connectivity limits, and in a vehicle, a perception or planning failure can require intervention before a remote service or human operator can respond. Vehicle safety therefore depends on the conditions the system can handle locally and the fallback available when those conditions no longer hold.

Even when sensors detect an obstacle, unstable predictions across frames can defeat downstream tracking and human fallback. In March 2018 in Tempe, Arizona, an Uber self-driving test vehicle struck and killed a pedestrian (National Transportation Safety Board 2019). The system detected her 5.6 seconds before impact, but its classification of her was unstable. Repeated reclassification among "vehicle," "other," and "bicycle" caused path prediction to discard tracking history. The final investigation placed recognition of the imminent collision at 1.2 seconds before impact. The system then suppressed action for one second without alerting the operator; only afterward did it issue

↩ The 2016 crash, read as an edge model whose fallback was a supervising driver, is established in section 11.11.4.

an alert and plan gradual slowdown. The operator did not intervene in time. The case connects perception, tracking, response timing, and human fallback, and it shows that detecting an object did not establish that the system would act safely on that detection.

An inspection drone surveying power lines may likewise lose reliable localization when lighting changes or repetitive textures confuse visual odometry. An alternate inference or control path provides **fallback redundancy** when the primary path cannot support the required decision, but its value depends on independent failure coverage and response time. A local sensor-based controller or validated landing mode may be more useful than a remote model that cannot answer before the control deadline. Edge AI must provision that response within the device's resource limits; a data center's additional replicas would not by themselves correct the same perception error.

Embedded systems can operate through critical intervals in which external intervention is too slow. Remote maintenance may remain possible between those intervals, but an autonomous landing or control maneuver needs a local response before its deadline. Validation must therefore cover both the component's failure modes and the system behavior that follows, including any ML model in the decision loop. NASA's Mars Polar Lander was lost in 1999 when, according to its review board, a spurious touchdown indication likely shut down the engines before landing (NASA Mars Program Independent Assessment Team 2000). For a learned component in such a control loop, validation must connect plausible input and sensor failures to the actions they can trigger, because exhaustive testing of every possible input is not a feasible safety argument.

Commercial aviation shows the same inherited hazard. A 2015 FAA airworthiness directive followed Boeing's discovery that a 787 powered continuously for 248 days could lose all AC power if all four generator control units entered failsafe mode at once,[3] so uptime itself became the risk factor (Federal Aviation Administration 2015). An endurance test shorter than that interval would not have exposed the defect. A learned component inherits the same hazard whenever deployment reaches a condition its evaluation never covered, such as accumulated state, a rare season, or an unusual sensor reading. Safety-critical systems[4] impose stringent reliability requirements precisely because such latent hazards surface only in the field. Across these settings, robustness depends on connecting the failure signal to a bounded response, and the fleet framework must distinguish the source of that signal before choosing a defense.

[3] **Failsafe mechanism**: A mechanism that moves the system to a validated safe state when a fault is detected. Circuit breakers, model fallback, rejection, and rollback are possible components, but none is inherently safe for every application. The service must establish that its chosen response meets the hazard and timing constraints; rejection thresholds also trade accepted workload against residual error and review capacity.

[4] **ASIL (automotive safety integrity levels)**: ISO 26262 uses automotive safety integrity levels to scale functional-safety obligations according to the hazard analysis. ASIL D is the most stringent level, but it is not a universal component-reliability percentage. An ML perception component complicates the resulting safety case because its behavior depends on the operational distribution and cannot be justified by aggregate test accuracy alone.

14.3 A Unified Framework for Robust AI

As those incidents show, a low-level fault and an environmental change can produce the same silent degradation. One flipped exponent bit in accelerator memory scales the stored value by a power of two and can turn a trained weight of 0.042 into a finite value near 1.4×10^{37} that passes every NaN check (section 14.9.1), while a gradual shift in user demographics erodes recommendation quality with latency and uptime flat. Production ML systems cannot treat these as isolated bugs. A unified framework must map how low-level hardware faults, software bugs, data drift, and adversarial inputs propagate upward into the model's output, starting with how distributed state sharing amplifies a single local failure across the fleet.

14.3.1 Shared state and fault surfaces across the fleet

The mechanisms that make a fleet efficient also share state or work across many machines. In distributed training, a single node transmitting corrupted gradients during an AllReduce operation can therefore poison the global model weights across every worker. Fault detection and recovery mechanisms keep training jobs running and support inference-time availability, yet a restored checkpoint or a responsive replica does not establish that the model behaves acceptably on current inputs. Robustness extends the recovery argument to graceful degradation, in which a reduced service must still meet a validated behavioral contract when inputs are adversarial or components degrade. Containing shared training state requires evidence that corrupted updates are detected or isolated. Byzantine fault tolerance protocols address the case in which a contributor sends arbitrary or malicious values rather than simply stopping. Ordinary hardware corruption also demands the fault diagnosis and recovery mechanisms already established for the fleet.

Serving repeats the same sharing at inference time. Batching, request routing, and pipeline-parallel stages place many users' requests on one model replica or one stage, so a corrupted replica or stage returns wrong answers to every request routed through it while its health checks still pass.

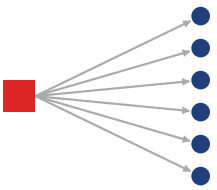

One corrupted gradient in AllReduce reaches every replica's weights.

↩ Gradient synchronization through AllReduce, in which every replica receives the same reduced gradient, is established in section 6.3.1.3.

↩ Byzantine failures, in which a component returns arbitrary values rather than stopping, are established in section 7.3.3.

When the cause is deliberate rather than accidental, the threat-vector classification in section 13.4.1.2, by attacker access, knowledge, and timing, defines the attacker that a defense must assume.

Large dense models amplify these risks. A GPT-3-class dense transformer, standing in here for the Archetype A (GPT-4/Llama-3) reference model, has 175B parameters and needs 350 GB for its weights alone at FP16/BF16 serving precision ($s_{\text{elem}} = 2$ bytes/param, with no optimizer state). Because that footprint exceeds the 80 GB per-device memory capacity M_{device} of a single A100- or H100-class accelerator, keeping these weights accelerator-resident requires distributing them across multiple devices. Each additional pipeline or tensor-parallel stage increases the fault surface compared with a monolithic deployment. A single bit flip or network partition in one stage can fail the entire inference request, and silent corruption can propagate through downstream computations that consume the affected values. Whether it changes the final prediction depends on the corrupted value, the computation, and the available detection or containment.

↩ Batching, request routing, and pipeline-parallel serving are established in chapter 10.

Capacity pressure also pushes teams to compress the model. Efficiency techniques such as INT8 quantization, pruning, and knowledge distillation let a model fit on fewer devices or serve more requests, but they change the deployed computation itself, so each one must be checked against the robustness evidence gathered for the uncompressed model. The quantity at stake is the model's **robustness margin**, the amount of input perturbation, numerical error, or representation change the model can absorb before its prediction changes. Each of these techniques can shrink that headroom, so a compressed model that matches the original's clean-data accuracy has not thereby shown that its margin survived. An optimized system therefore needs reliability-engineering evidence[5] that its error detection and recovery mechanisms cover the faults in its operating model and that the resulting predictions still meet the service contract.

Scale also compounds the chance that silent corruption occurs somewhere in the fleet. Under the fleet model of section 7.5.2, N devices, each with independent residual-error probability p over the same interval, see at least one silent error with probability $\Pr(\geq 1) = 1 - (1-p)^N$. A per-device rate that is negligible on one machine therefore becomes a routine event on a fleet of tens of thousands of accelerators, which turns a per-device reliability number into a coordination cost, the third axis of C^3. Once a silent error is likely somewhere in the fleet, operators must agree on which replicas, checkpoints, and outputs remain trustworthy before they can contain it, and that agreement requires evidence that ordinary health checks do not collect. Hardware faults, adversarial inputs, and software defects all create this evidence problem, which is why the chapter treats them within one framework.

5 **Reliability engineering**: Uses fault models and failure-mode analysis to connect component behavior to system outcomes. A series model gives system reliability as $R_{\text{system}}(t) = e^{-N_{\text{series}}\lambda t}$ when all N_{series} components must survive, failures are independent, and each has the same constant exponential failure rate. ML deployments also need behavioral evidence for drift, crafted inputs, and uncertainty; the component-survival model alone does not cover these failure modes.

14.3.2 The three pillars of robust AI

When a model's error rate climbs without a service outage, an engineer must determine whether customer behavior shifted, an attacker perturbed the inputs, or an accelerator corrupted a weight before choosing a defense. Each of the three pillars named at the start of the chapter decomposes into concrete failure subtypes, and each subtype pairs with the defense that targets it (figure 14.2).

The subtype names matter less than the evidence each pillar produces, because the evidence is what an operator actually observes and what selects the defense family. Environmental shifts can produce statistical differences between production data and a reference population. Those differences motivate investigation, but an input-distribution monitor cannot detect every change in the relationship between inputs and outcomes. Monitoring, recalibration, and retraining address different parts of that problem; outcome evidence helps determine which response is warranted.

Input-level attacks deliberately target model behavior through crafted inputs or poisoned training samples. Anomalous inputs can raise suspicion, but an effective attack need not expose a recognizable perturbation or statistical outlier. Adversarial training, certification, and input sanitization address different threat models.

System-level faults originate in the hardware, code, frameworks, and deployment infrastructure that support ML systems: numerical instability, preprocessing bugs, race conditions and replica desynchronization, dependency mismatches, and hardware faults such as bit flips. Many produce explicit errors or outages, but a fault that leaves the service answering appears instead as changed model behavior that can pass for either of the other pillars. Section 14.9 develops the diagnostic signals that separate those masquerades after the first two pillars have been treated. All three pillars share the same operational control loop of detection, degradation, and adaptation.

↩ The mechanics of hardware and software faults, and their detection and recovery, are established in chapter 7.

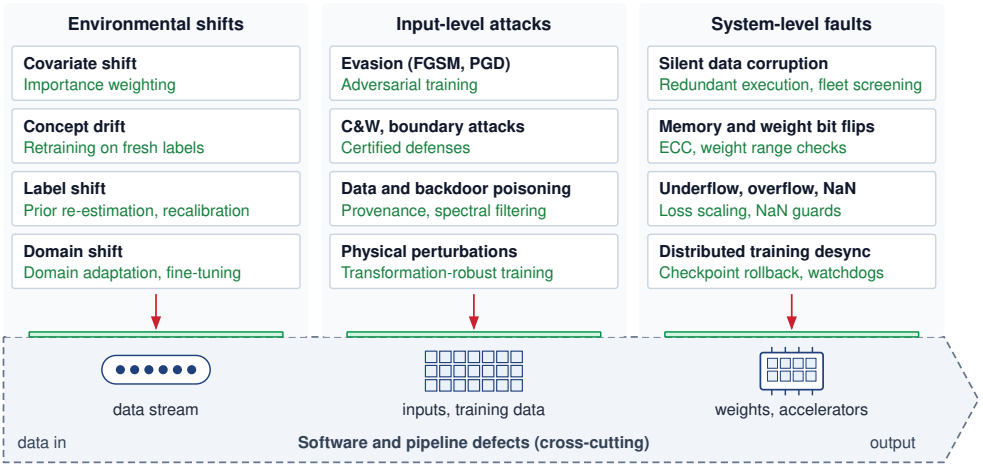

Figure 14.2: The Robustness Taxonomy: Environmental shifts, input-level attacks, and system-level faults decomposed into concrete failure subtypes (dark), each paired with its principal defense (green). Each pillar descends onto a defense wall at the point where it enters the ML pipeline, the data stream, the model's inputs and training data, or the weights and accelerators. Software and pipeline defects originate in the system layer, but they are drawn as the cross-cutting pipeline rail on which all three walls sit because their symptoms can amplify or imitate any of the pillars.

14.3.3 Common robustness principles

An uncalibrated monitor that fires on harmless traffic fluctuations can trigger unnecessary retraining or flood a fallback service just as easily as an undetected failure can corrupt predictions. Across all three categories, the shared engineering problem is deciding which signal triggers which response. Robust systems need a detection threshold, a degradation path, and an adaptation mechanism, each with an explicit cost budget.

Detection and monitoring form the foundation of that strategy, and each pillar asks for a different signal because each failure originates in a different layer. System-level monitoring samples hardware and runtime metrics for evidence of temperature anomalies, voltage fluctuations, memory errors, or silent data corruption; detection can limit further propagation, although the first observable symptom may follow a corrupted state update. Where system monitors watch the machinery, input-level attack monitoring uses statistical or activation-space tests to flag suspected adversarial inputs and poisoning attempts for rejection or investigation. Environmental-shift monitoring compares production traffic with reference distributions using tools such as maximum mean discrepancy (MMD),[6] population stability index (PSI), or Kolmogorov-Smirnov (KS) tests, the statistics developed in section 14.4.2. Each of these monitors also has a cost. Robustness mechanisms must be budgeted, not merely enabled, because every detector trades sensitivity against false positives, latency, and compute overhead.

Graceful degradation turns detection into a bounded operating mode instead of a crash, stepping down a predeclared ladder of cheaper models and simpler features rather than failing outright. Robustness adds a behavioral condition to that ladder, because the degraded model and its fallback must still satisfy the service's acceptance criteria under the failure that triggered them. A quantized replica that keeps serving when capacity drops therefore needs its own validation under the degraded operating conditions, since quantization can shrink the robustness margin even when clean accuracy holds. Ensemble fallback systems likewise preserve service only if the alternate models cover the relevant failure and complete within the serving deadline. Their quality under that failure, switching delay, and capacity under diverted traffic require evaluation, and ensemble membership alone establishes none of these properties.

Adaptive response completes the loop by changing system behavior when the signal persists, for example by activating error correction, adding input preprocessing, or updating model parameters, so robustness is an ongoing adjustment rather than a one-time configuration. Across the three pillars

6 **Maximum mean discrepancy (MMD)**: A discrepancy between kernel mean representations of distributions, used to construct two-sample tests. A kernel on input vectors lets MMD respond to joint-distribution changes that per-feature comparisons can miss; sensitivity depends on the kernel and test calibration. Direct pairwise estimation costs $\mathcal{O}(n^2)$ in sample size, while linear-time estimators and approximations trade computation against statistical efficiency.

↩ Graceful degradation as a ladder of model and feature fallbacks, each validated before it is needed, is established in section 7.12.

of figure 14.2, error-correcting code (ECC) memory and redundant execution detect faults, weight range checks and checkpoint rollbacks contain the damage, and retraining or adversarial training adapts the model. Each pillar therefore needs its own version of all three mechanisms, and no single defense covers the taxonomy.

The evaluation plan must preserve those distinctions, and it needs one row that is not a pillar at all. Population monitors ask whether traffic as a whole has changed, while a serving system must also decide, request by request, whether this particular input lies inside the model's demonstrated competence. The tools for that per-request decision, namely calibration, out-of-distribution (OOD) scores, and selective prediction with its risk-coverage curve, are developed in section 14.4.4, after the drift statistics they complement. Each failure class needs its own minimum evidence (table 14.1), so a passed test in one row cannot become a robustness claim about another. Per-request checks separate near-OOD inputs, which resemble the training data but fall outside its classes or conditions, from far-OOD inputs drawn from unrelated sources. Adversarial evaluation measures robust accuracy under a stated attack and certified accuracy where a proof is available (section 14.6.1.1), and poisoning evaluation tests contributors that use fake identities (sybils) or send arbitrary updates (Byzantine contributors) (section 14.7).

Table 14.1: Robustness Evaluation Matrix: Each failure class requires its own intervention, evidence, slices, and operational decision. Aggregate test accuracy or one attack benchmark cannot stand in for the complete matrix.

Failure class	Evaluation set or intervention	Primary evidence	Required slice	Decision boundary
Natural shift	Time-, geography-, device-, and environment-held-out data	Accuracy or loss, calibration error, and drift statistics	Critical cohorts and tail conditions	Continue, investigate, recalibrate, retrain, or fall back
Per-request competence	In-distribution errors plus near- and far-OOD inputs	Risk-coverage curve, OOD score, and fallback load	High-consequence request classes	Answer, route, abstain, or reject
Adversarial manipulation	White-box, black-box, transfer, and physical attacks matched to the threat model	Robust accuracy by attack budget; certified accuracy where available	Targeted classes and attacker capabilities	Accept, harden, restrict, or block release
Poisoning	Source-compromise, label-flip, backdoor, fake-identity (sybil), and Byzantine-contributor scenarios	Clean and attack-success metrics, slice loss, and detector recall	Sources, contributors, and subpopulations	Quarantine, exclude, reweight, or roll back
System fault	Fault injection across data, runtime, accelerator, and network paths	Silent-error rate, detection latency, and recovery success	Correlated and multi-component faults	Retry, isolate, degrade, or fail safe

14.3.4 Integration across the ML pipeline

A serving-time filter cannot repair a fraud model that was trained on poisoned records, and an ingestion filter cannot stop an attacker who perturbs inputs at inference time. Robustness therefore requires evidence across the entire ML lifecycle, with complementary controls at each stage applying the defense-in-depth principle to model behavior. Ingestion checks flag malformed or statistically anomalous records before training, training-time defenses target specified perturbations, validation stress-tests behavior on out-of-distribution data within a named deployment envelope, and serving-time filters and monitors watch live traffic. Passing one stage's check says nothing about the failures another stage must test.

A production fraud detection pipeline therefore evaluates its ingestion checks against poisoned training records and its serving-time filters against evasion attempts separately, and neither guarantees interception. The monitoring layer acts as the safety net, detecting distribution drift, such as a sudden shift in transaction amounts or user geolocations, and opening an investigation. Retraining follows only when validated outcome evidence, root-cause analysis, and the release gate show that a model update is the appropriate response.

↩ Defense in depth, which requires independent layers so that defeating one does not defeat the system, is established in section 13.4.1.3.

Hardware faults (transient, permanent, and intermittent) and the fleet's mechanisms for detecting and recovering from them are classified in section 7.4. Robustness evaluation adds a behavioral test to those mechanisms: after a fault is detected or a request retried, the output must still satisfy

the model's acceptance criteria. A retry alone does not establish that the result is correct. The same distinction between service health and prediction quality applies when the machinery works correctly but the world evolves beyond its training data.

Checkpoint 14.1: Diagnosing the failure signal

The unified framework requires identifying which of the three pillars produced a silent failure before choosing a defense, and recognizing that software faults can masquerade as any of them.

Classifying the threat

- ☐ **Pillar classification**: classify a model losing accuracy as customer behavior evolves, a crafted input that forces misclassification, and an accelerator HBM bit flip under the correct robustness pillars.
- ☐ **Software-fault mimicry**: explain why a preprocessing bug that shifts every inference's feature statistics is hard to distinguish from genuine environmental shift, then name the pillar it belongs to by origin and the pillar whose symptom it imitates.

Choosing the response

- ☐ **Response selection**: map an ECC-corrected single-bit memory error and a fired drift monitor to the shared responses of detection, graceful degradation, and adaptive response.
- ☐ **Fleet-scale effects**: explain why a negligible per-device silent-corruption rate becomes an architectural constraint at a fleet scale of tens of thousands of accelerators.

14.4 Environmental Shifts

A credit-scoring or recommendation model trained on last quarter's user behavior can lose accuracy as seasonal demand, demographics, or economic conditions move, because training data freezes a past world while production traffic keeps changing. Environmental shifts arise whenever data distributions, user behavior, or deployment contexts depart from the training population, and a shift becomes a robustness failure when model behavior no longer meets the service contract. Changed conditions can also expose latent software defects or alter the effectiveness of an adversarial defense, so choosing the right intervention requires first separating which statistical component of the data-generating process has moved.

14.4.1 Distribution shift and concept drift

That separation is already necessary when a model moves across sites without any change in the underlying task. A medical diagnosis model trained on X-ray images from one hospital and deployed in a clinic with different imaging equipment encounters altered contrast and noise even for the same medical conditions. This difference between training and deployment populations is **distribution shift**; whether it reduces diagnostic accuracy requires evaluation on the clinic's data. A move to a different source altogether, such as a new hospital, sensor, or region, is often called domain shift, and it can combine several of the shift types distinguished below.

Distribution shift takes three geometric and statistical forms (figure 14.3). Inputs can migrate to an unfamiliar region of feature space (panel a), label priors can skew (panel b), or the underlying mapping itself can change (panel c). Here, x denotes an input and y its outcome label. Input-only monitors can detect changes in $p(x)$; distinguishing changes in label frequencies or in $p(y \mid x)$ generally requires outcome evidence or additional assumptions.

14.4.1.1 Technical categories

A vehicle perception model trained mostly on daytime images encounters unfamiliar pixel statistics at night, or in rain, snow, and fog, even when the physical objects corresponding to each label have not changed. That form of shift is **covariate shift**, which changes the input distribution while the relationship between inputs and outputs stays fixed (Quiñonero-Candela et al. 2009). Treating a deployment change as covariate shift requires assuming that the conditional labeling relationship stays fixed; an unchanged task name alone does not establish that assumption. Any accuracy loss depends on the luminance shift, the sensor, and how much adverse-condition data the training set contained, so each deployment must measure it on held-out night and weather slices rather than assume a published figure. Under the covariate-shift assumption, the changed population can place more probability in regions the model learned poorly, as illustrated in panel (a) of figure 14.3. Aggregate accuracy can then decline even though the labeling rule remains fixed. This is a change in the distribution of errors, not a claim that every affected request becomes incorrect. Because the labeling rule is fixed, the remedy can reweight existing labeled data rather than relabel it. Importance weighting gives each training example a weight equal to how much more or less likely its input is under deployment traffic than under training traffic, so the weighted training loss estimates the loss on deployment traffic. That estimate holds only where the training data cover the new inputs, because no weight can recover a region that contains no training examples.

Even a shift that was designed to be no shift at all costs accuracy. Recht et al. (2019) rebuilt the ImageNet validation set a decade after the original by following the original collection and labeling protocol, then evaluated existing models on the new images. In the larger testbed of Taori et al. (2020), every one of the 212 models that reach at least 55 percent on the original set scores lower on

↩ A fleet drift monitor that must tell input-distribution change from input-output change is established in section 12.6.4.3; the covariate, label, and concept shifts it separates are defined here.

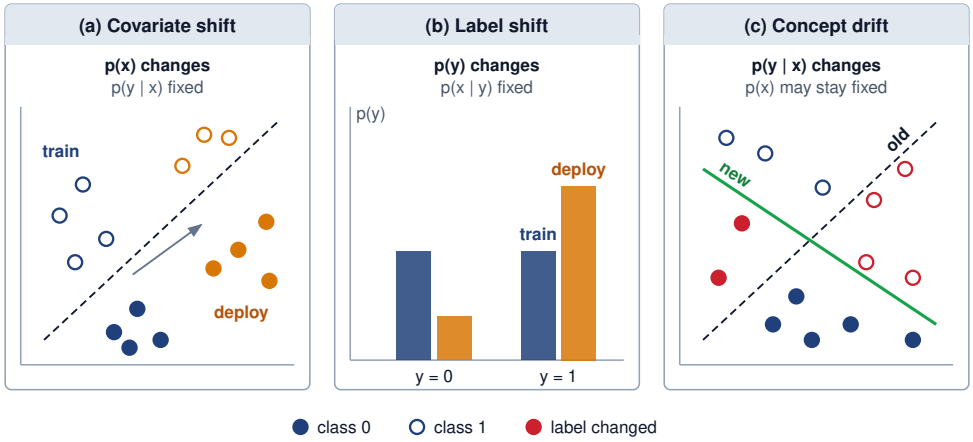

Figure 14.3: Types of Distribution Shift: (a) Under covariate shift, deployment inputs move to a new region of feature space while the labeling rule, and so the correct boundary, stays fixed. (b) Under label shift, class frequencies change while each class's inputs look the same; the bar heights are illustrative. (c) Under concept drift, the same inputs acquire new labels, so points between the old and new boundaries become errors. Covariate shift changes $p(x)$ with $p(y \mid x)$ fixed, label shift changes $p(y)$ with $p(x \mid y)$ fixed, and concept drift changes $p(y \mid x)$; the distinction determines whether reweighting, recalibration, or model adaptation is appropriate.

the new one, by 7 to 14 points with a median of 12 (figure 14.4). The loss follows a nearly straight line, so a model that is better on the original set is better on the new one, yet no architecture, training recipe, or amount of extra data closes the gap. Whatever separates the two sets is not an obvious visual change, yet it is costly to every model. For a fleet, accuracy measured on held-out data from the training distribution is therefore an optimistic estimate of accuracy on the next population of inputs, not a prediction of it.

Covariate shift moves the inputs while the true mapping $p(y \mid x)$ stays fixed. A more insidious variant occurs when the mapping itself changes, so that the correct label for a given input today differs from its label during training.

Definition 14.2: Concept drift

Concept drift is the deployed-model subtype of distribution shift in which the statistical relationship $p(y \mid x)$ changes over time, potentially changing the correct class decision or its probability even when the input distribution stays fixed. Data drift names a change in $p(x)$; the more restrictive covariate-shift case additionally assumes that $p(y \mid x)$ remains stable.

1. **Significance**: It causes silent model degradation because the historical mapping learned by the model is no longer representative of current reality. It also shortens the time a model can serve before retraining is required, so fraud-detection models, recommender systems, and other behavior-dependent models may need periodic retraining or recalibration as adversaries, users, and policies change. The remediation cost can range from recalibration or incremental fine-tuning to a full training run. Faster drift can raise the amortized cost per prediction because remediation recurs more often, although not every remediation costs as much as the original training run.
2. **Distinction**: Under covariate shift, $p(x)$ changes while $p(y \mid x)$ remains stable, so resampling or reweighting may help where the reference data cover the new inputs. Detecting data drift alone does not establish that the conditional relationship stayed fixed. Concept drift instead requires adapting to a new $p(y \mid x)$, usually using recent ground-truth outcomes. This can make concept drift structurally more expensive to remediate than resampling an existing labeled distribution.

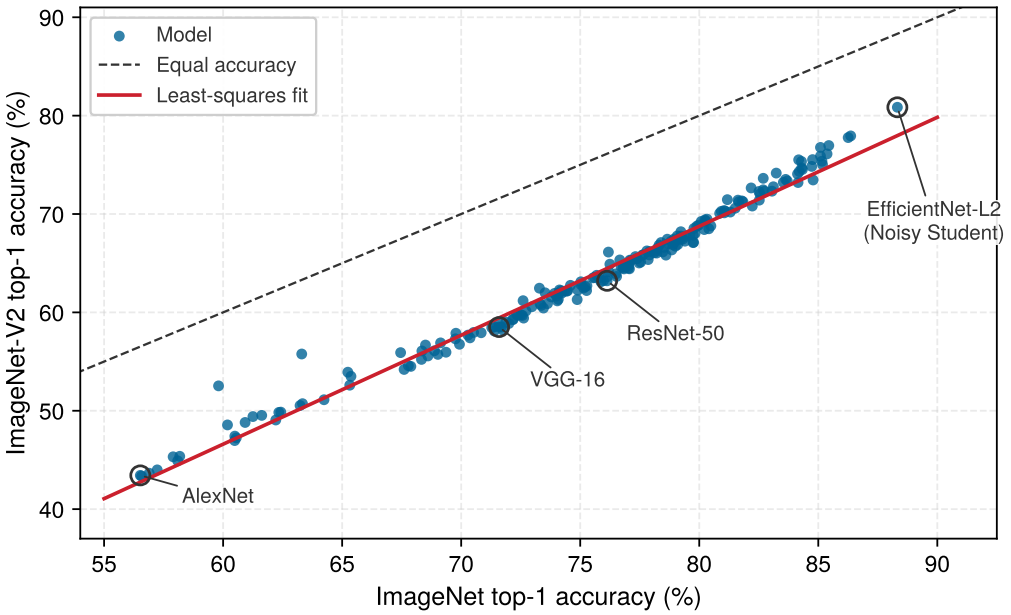

Figure 14.4: Accuracy on a Fresh Test Set from the Same Recipe: Top-1 accuracy on ImageNet-V2, a test set rebuilt by following the original ImageNet collection protocol, against accuracy on the original ImageNet validation set, for testbed models that reach at least 55 percent on the original. Every point lies below the dashed line of equal accuracy, and the red least-squares line shows that the loss tracks accuracy almost linearly. ImageNet-V2 is from Recht et al. (2019); model evaluations from the testbed of Taori et al. (2020).

3. **Common pitfall**: A frequent misconception is that concept drift is detectable by monitoring input feature statistics. Because $p(x)$ may be entirely unchanged, input-level monitoring (PSI, KL divergence on features) will show no signal. An outcome-based detector compares predictions with ground-truth outcomes and must wait for those labels to confirm a change. Proxy signals can still justify investigation or containment before that confirmation.

Because input monitors can miss concept drift entirely, classic drift-adaptive learners watch error on recently labeled examples and shrink or grow their training window when that error changes, rather than tracking feature statistics (Widmer and Kubat 1996).

A production detector that relies on labeled errors must wait for those outcomes. Input statistics and other proxy signals may justify earlier investigation or containment, but they do not by themselves establish that the input-output relationship has changed.

Whereas concept drift alters the conditional mapping for a given input, a hospital's case mix during an epidemic such as the COVID-19 pandemic can change the prevalence of respiratory illness while the radiographic appearance within each diagnosis remains stable. This form is **label shift**, in which the label marginal $p(y)$ changes while the class-conditional input distribution $p(x \mid y)$ remains fixed (Lipton et al. 2018). By Bayes' rule, the posterior probability $p(y \mid x)$ is proportional to $p(x \mid y)p(y)$, so the probability of a diagnosis given an observed image changes whenever prevalence $p(y)$ changes, even though each class looks the same inside the scanner. A changed posterior therefore does not by itself indicate concept drift, and the three shifts are best separated by what stays fixed. Covariate shift keeps the input-to-label rule fixed, label shift keeps each class's input distribution fixed, and concept drift changes the input-to-label rule in a way that correcting for the new class frequencies cannot undo, because the inputs typical of each class have themselves changed.

Because the within-class feature distribution $p(x \mid y)$ is unchanged, a fixed classifier's confusion matrix on held-out validation data also remains valid within each true class, which allows operators to estimate the new deployment prevalence from unlabeled predictions. That matrix records, for

each true class, how often the classifier predicts each class. On deployment traffic, the observed frequency of each predicted class then matches, up to sampling noise, the matrix applied to the unknown new class frequencies, so inverting the confusion matrix recovers the new prevalence. Multiplying each posterior probability by the ratio of new to old prevalence for its class, then renormalizing, recalibrates the model without retraining. A real clinical population may violate the stable within-class assumption, so prevalence change alone does not establish pure label shift, and confirming that assumption requires outcome labels.

Unlike adversarial attacks that require malicious intent, these three forms of shift emerge organically as language evolves with new slang, sensors age, or market conditions move. When a deployed model actively influences user behavior, as in fraud detection, recommendation engines, or credit underwriting, the shift is not merely exogenous. The sociotechnical feedback invariant (principle 27) governs these performative environments, in which the model's own predictions change the behavior it later observes. In them, the probability distribution of future inputs $p_{t+1}(x)$ depends directly on the model's past decisions $f_t(x)$ through the environment's response function $g(p_t(x), f_t(x))$. Deployed predictions can induce user and adversary adaptations that accelerate distribution shift, compressing the model's viable operating window and demanding closed-loop monitoring rather than static validation. Section 16.9.4.1 returns to these feedback loops.

The taxonomy names what changed; it does not explain why a change that leaves the labeling rule intact can still hurt. Often the reason is a shortcut the model learned from its training data, which a covariate shift then exposes. A classic example is a model that learns to identify "cow" by detecting "grass" background. When presented with a cow on a sandy beach, the model fails. The underlying cause is a **spurious correlation**: a feature that is predictive in the training set but not causally related to the label.

Standard training by empirical risk minimization can favor a shortcut that predicts labels well in the training distribution, even when that association fails in deployment. Background, texture, or shape may provide useful evidence, but any of them can become a shortcut when its training association does not hold across contexts. Techniques like **group distributionally robust optimization** explicitly mitigate this by minimizing the *worst-case group loss* (for example, cows on sand) rather than the *average loss*. The method requires groups to be known or inferred in advance, and optimizes the worst represented-group loss rather than guaranteeing performance in every context. Its effect must be evaluated on the named groups and on held-out contexts; unseen groups remain outside that objective's direct coverage, so deployed systems still need runtime drift metrics to detect when live traffic leaves the validated distribution.

14.4.2 Measuring distribution drift

An operations team that retrains on every visible feature shift wastes compute on harmless fluctuations, while a team that waits for customer complaints misses silent degradation in a critical cohort. Deciding whether a suspected shift matters requires a measured signal rather than an impression. Drift metrics must answer an operational question: whether the model should keep serving, be monitored more closely, or enter a root-cause investigation that may justify recalibration, data repair, or retraining. PSI supplies a cheap fleet-wide alert that starts that decision, while density divergences and hypothesis tests let a service calibrate thresholds without mistaking them for universal policy.

14.4.2.1 Population stability index (PSI)

When a fleet-wide drift alert fires, an investigation must identify which feature range or category moved rather than rely on a single aggregate score. PSI supports that search because it is built bin by bin. A continuous feature is first cut into value ranges and a categorical feature keeps its categories, and equation 12.10 defines PSI as a sum over the resulting n_{bins} bins. For each bin it multiplies the difference between the current bucket proportion A_b and the baseline proportion E_b by the logarithm of their ratio. The same sum therefore decomposes directly into each bin's contribution. Bin b contributes $(A_b - E_b)\ln(A_b/E_b)$, and each such term is nonnegative, because $A_b - E_b$ and $\ln(A_b/E_b)$ always share a sign. Swapping the two proportions flips the sign of both factors and leaves each bin's term unchanged, so PSI is symmetric in the baseline and current distributions. The logarithmic factor penalizes large relative changes, while the difference weights them by absolute magnitude, so a per-bin decomposition shows which feature ranges or categories contribute to an

alert. Identifying an affected population cohort requires connecting those bins to the corresponding requests.

The fleet convention bands of section 12.6.4.3 map PSI values to negligible, moderate, and major shift. They are monitoring conventions, especially common in credit-scoring practice, rather than universal statistical guarantees, so a service validates them against its own binning, baseline, sample size, predictive features, and response cost, and reads PSI together with feature importance and downstream model-performance metrics (Yurdakul and Naranjo 2020).

Binning and the treatment of empty bins change what PSI can see, so the bin boundaries chosen on the reference sample must be reused for the comparison population. An empty bin is handled with a smoothing constant ϵ_{smooth} and renormalization; because that choice can materially change PSI on sparse bins, the monitoring configuration records it alongside the reference.

Monitoring population stability index over time as plotted in figure 14.5 reveals when the input population moves from the stable band into the warning band and then the major-shift band. That movement triggers an escalation path that may lead to retraining after performance correlation. At week 41, the service approves and versions a new input reference after validating the changed operating population. The lower PSI thereafter compares traffic with that new reference, not the original one. Changing model weights alone would not reduce input PSI against a fixed reference, and the discontinuity does not establish improved predictions.

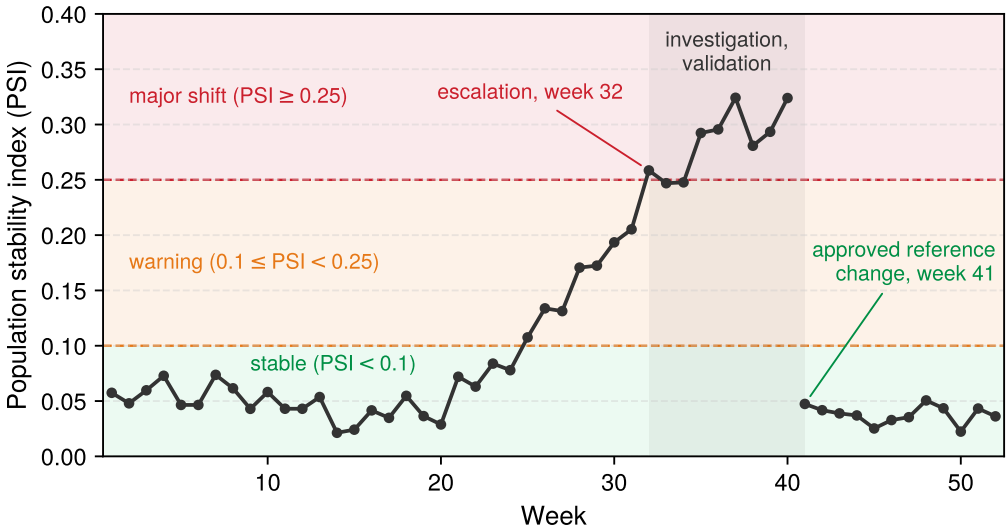

Figure 14.5: A Calibrated Distribution-Shift Workflow: In this synthetic service, PSI enters the warning zone as user behavior shifts and crosses the service's escalation threshold at week 32. Following investigation and validation of the changed operating population, the service versions a new input reference at week 41. Values on either side of that reference change use different comparison populations; the drop is not evidence of model recovery.

This threshold-based path automates alerting, not remediation. Operators still need labels or other outcome evidence, root-cause analysis, and an evaluated response before changing the model. Forcing continuous variables into discrete bins also discards information about distribution shape, motivating continuous drift metrics.

14.4.2.2 Kullback-Leibler divergence

For continuous features, the Kullback-Leibler (KL) divergence compares fitted densities rather than fixed bins, at the price of the density estimate's own assumptions and error. From baseline distribution P_0 to current distribution P_t it is:

$$\mathcal{D}_{\text{KL}}(P_0 \| P_t) = \int_{-\infty}^{\infty} P_0(x_{\text{feat}}) \ln\left(\frac{P_0(x_{\text{feat}})}{P_t(x_{\text{feat}})}\right) dx_{\text{feat}}$$

where $P_0(x_{\text{feat}})$ and $P_t(x_{\text{feat}})$ are the probability density functions of the baseline and current distributions, respectively. Unlike PSI, KL divergence is asymmetric: $\mathcal{D}_{\text{KL}}(P_0 \| P_t) \neq \mathcal{D}_{\text{KL}}(P_t \| P_0)$. The fraud-monitoring example in section 14.4.2.4 uses $\mathcal{D}_{\text{KL}}(\text{baseline} \| \text{current})$, measuring how much information is lost when using the current distribution to approximate the baseline.

To address asymmetry, practitioners often use the Jensen-Shannon divergence:

$$\mathcal{D}_{\text{JS}}(P_0 \| P_t) = \frac{1}{2}\mathcal{D}_{\text{KL}}(P_0 \| \bar{P}) + \frac{1}{2}\mathcal{D}_{\text{KL}}(P_t \| \bar{P})$$

where $\bar{P} = \frac{1}{2}(P_0 + P_t)$ is the mixture distribution. **Jensen-Shannon divergence** is bounded between 0 and $\ln(2)$ (approximately 0.693), making threshold selection more intuitive than unbounded KL divergence.

KL divergence has no conventional bands comparable to PSI's, so each service sets its own. The fraud service analyzed in section 14.4.2.4 starts from the bands in table 14.2 and validates them against its historical traffic.

Table 14.2: Scenario-Calibrated KL Alert Bands: The worked fraud service uses these values to assign monitoring severity, not to command retraining. Another feature, density estimator, or consequence model requires different calibration.

$\mathcal{D}_{\text{KL}}$ value	Interpretation
$\mathcal{D}_{\text{KL}} < 0.05$	Minimal divergence
$0.05 \leq \mathcal{D}_{\text{KL}} < 0.1$	Moderate divergence
$\mathcal{D}_{\text{KL}} \geq 0.1$	Significant divergence

For practical computation, kernel density estimation with Gaussian kernels provides smooth density approximations suitable for integration. Direct evaluation at every sample against every kernel costs $\mathcal{O}(n^2)$ for n samples. Sampling, a fixed evaluation grid, or faster approximation methods change this cost; bandwidth selection also changes the estimated divergence and must be part of threshold calibration.

14.4.2.3 Statistical significance testing

PSI and KL divergence quantify how large a distributional change appears to be; statistical hypothesis tests ask the complementary question of whether the observed difference is larger than sampling noise. The simplest case is a shift in the mean of one feature whose spread is known, which a normal approximation can test directly.

> Napkin Math 14.1: Detecting a real distribution shift
>
> **Problem**: A model monitors a critical input feature with a known standard deviation of 0.3. The baseline mean was 0.5. Over the last 1,000 requests, the mean has shifted to 0.55. The engineering question is whether this is a random fluctuation or a real distribution shift.
>
> **Math**: The calculation treats the baseline mean as known and the requests as independent observations with the stated standard deviation. A normal approximation then measures how unusual the observed mean would be under that baseline. If the baseline mean is itself estimated from a finite sample, its uncertainty also contributes to the comparison.
>
> 1. **Difference in means**: 0.05.
> 2. **Standard error**: $0.3/\sqrt{1,000} \approx 0.009$.
> 3. **Statistical significance**: The shift is approximately 5.3 standard errors away from the mean.
> 4. **P-value**: < 0.001.
>
> **Systems insight**: Statistical significance assesses the observed change against sampling variability under the baseline model. A shift of 0.05 might seem "small," but under the baseline model, the probability of observing a sample mean at least this far from the baseline is less than

0.1 percent. In the machine learning fleet, this supports a drift alert under the test assumptions, not yet a confirmed model regression. The system should trigger investigation, increased monitoring, and correlation with precision, recall, latency, and business metrics. Feature importance and shift magnitude help prioritize that investigation; fallback or retraining requires evidence that the response addresses the cause or contains the resulting risk.

A mean test sees only a change in location and needs the spread in advance. The two-sample Kolmogorov-Smirnov (KS) test (Berger and Zhou 2014) compares the empirical cumulative distribution functions (CDFs) of two samples without assuming any specific parametric form. The test statistic is:

$$\mathcal{D}_{\text{KS}} = \sup_{x_{\text{feat}}} |F_n(x_{\text{feat}}) - G_m(x_{\text{feat}})|$$

where F_n and G_m are the empirical CDFs of samples of size n and m respectively. For independent samples from continuous distributions, the following asymptotic rule rejects the null hypothesis of no distributional difference. Small samples, ties, or dependent traffic require an appropriately calibrated test.

$$\mathcal{D}_{\text{KS}} > c(\alpha_{\text{sig}}) \sqrt{\frac{n+m}{nm}}$$

where $c(\alpha_{\text{sig}})$ depends on the significance level (for example, $c(0.05) \approx 1.36$). The KS statistic is most sensitive to differences near the center of the distribution, such as a location shift, and less sensitive to differences confined to the tails.

For categorical features, the **chi-square goodness-of-fit test** compares observed frequencies to expected frequencies under the baseline distribution:

$$\chi^2 = \sum_{i=1}^{k} \frac{(n_i^{\text{obs}} - n_i^{\text{exp}})^2}{n_i^{\text{exp}}}$$

where n_i^{obs} is the observed count in category i and n_i^{exp} is the expected count based on the baseline distribution. With independent observations and well-populated expected counts, χ^2 is compared with the chi-square critical value for $k-1$ degrees of freedom at significance level α_{sig}.

When monitoring many features simultaneously, the significance test must also account for repeated comparisons. Applying Bonferroni correction (dividing α_{sig} by the number of tests) controls the probability of at least one false rejection within the specified test family. False discovery rate procedures instead target the expected fraction of false discoveries among rejections under their own assumptions. For N_{tests} monitored features, one test each, at significance level $\alpha_{\text{sig}} = 0.05$, Bonferroni requires each test to achieve $p_{\text{val}} < 0.05/N_{\text{tests}}$ for significance. This correction across features does not by itself control false alarms over an indefinitely repeated sequence of monitoring windows.

14.4.2.4 Production fraud detection drift analysis

Consider a fraud detection model serving an e-commerce platform with two key input features: user country (categorical) and transaction amount (continuous). After six months in production, the operations team suspects distribution drift and must decide whether to retrain. From its historical baseline the team has already configured three investigation triggers. The first is a population stability index above 0.1 on any high-importance feature. The second is a KL divergence above 0.05 on a continuous feature. The last is a KS test p-value below 0.01, the per-test level left after a Bonferroni correction across its monitored features.

Step 1 categorical feature analysis The user-country feature has four named countries and an Other bucket, and decomposing its PSI bucket by bucket shows which part of the traffic moved (table 14.3). Each row compares the baseline share against the current share so that localized shifts do not hide inside the aggregate score.

Table 14.3: Country-Feature PSI Decomposition: Per-country contribution to the total PSI for the user country feature in a production fraud-detection model. The "Other" bucket drives most of the divergence (0.0470 of the 0.065 total), reflecting growth of non-USA/UK/Germany/France traffic over six months in production.

Country	Baseline (E_i)	Current (A_i)	$E_i - A_i$	$\ln(E_i/A_i)$	Contribution
USA	0.45	0.38	0.07	0.169	0.0118
UK	0.20	0.18	0.02	0.105	0.0021
Germany	0.15	0.14	0.01	0.069	0.0007
France	0.10	0.12	-0.02	-0.182	0.0036
Other	0.10	0.18	-0.08	-0.588	0.0470

Summing the per-country contributions gives $\text{PSI}_{\text{country}} = 0.0118 + 0.0021 + 0.0007 + 0.0036 + 0.0470 = 0.065$. The PSI of 0.065 remains inside this fraud service's calibrated baseline for user country, so the aggregate feature does not independently escalate the investigation, yet the decomposition still matters. The Other bucket nearly doubled its share of traffic, so the team checks performance within that slice rather than relying on the aggregate PSI. The changed share warrants investigation without establishing that those transactions are less accurately classified.

Step 2 continuous feature analysis For the transaction amount feature (log-transformed for approximate normality), fitting Gaussian distributions to the samples yields a baseline distribution of $\mu_{\text{base}} = 4.2$ and $\sigma_{\text{base}} = 1.1$ log-dollars, compared with a current distribution of $\mu_{\text{curr}} = 4.5$ and $\sigma_{\text{curr}} = 1.3$ log-dollars.

For the fitted Gaussian distributions, KL divergence has the following closed-form solution. It is exact for those density models; its fidelity to the transaction populations depends on the fit.

$$\mathcal{D}_{\text{KL}} = \ln \frac{\sigma_{\text{curr}}}{\sigma_{\text{base}}} + \frac{\sigma_{\text{base}}^2 + (\mu_{\text{base}} - \mu_{\text{curr}})^2}{2\sigma_{\text{curr}}^2} - \frac{1}{2}$$

The closed-form solution evaluates to $\mathcal{D}_{\text{KL}} = \ln \frac{1.3}{1.1} + \frac{1.30}{3.38} - 0.5 = 0.167 + 0.385 - 0.5 = 0.052$. The KL divergence of 0.052 indicates moderate drift, warranting further investigation but not immediate retraining.

Step 3 statistical significance via KS test The KS test on 10,000 baseline samples and 10,000 current samples of transaction amount returns $\mathcal{D}_{\text{KS}} = 0.089$. The service tests at the Bonferroni-corrected level $\alpha_{\text{sig}} = 0.01$ that the team configured. At this per-test level, the asymptotic coefficient is $c(\alpha_{\text{sig}}) = \sqrt{-\ln(\alpha_{\text{sig}}/2)/2}$, so the critical value is $c(0.01)\sqrt{\frac{n+m}{nm}} \approx 1.63 \times 0.0141 \approx 0.023$.

The observed statistic is several times that critical value. Tightening the level tenfold, to $\alpha_{\text{sig}} = 0.001$, raises the coefficient $c(\alpha_{\text{sig}})$ only modestly, because it grows with the square root of $\ln(2/\alpha_{\text{sig}})$, so the same statistic also rejects at that stricter level. The difference is therefore statistically significant $(p_{\text{val}} < 0.001)$. However, *statistical significance* alone does not mandate retraining; the *practical significance*, the size of the shift that the PSI and KL values measure against the configured triggers, supports investigation while monitoring continues, rather than immediate retraining.

Step 4 decision framework application Set against the team's configured triggers, the three results split into one signal to monitor and two to investigate (table 14.4). Comparing all three metrics side by side prevents a single statistical rejection from bypassing the practical-significance check.

Table 14.4: Fraud-Monitoring Decision Matrix: Combined quantitative evidence for this worked scenario, compared with thresholds calibrated to its historical traffic and review capacity. The KL and KS results cross their investigation triggers; serving continues with increased monitoring while the team assesses the shift. Statistical significance alone does not command retraining.

Metric	Value	Threshold	Action level
PSI (country)	0.065	< 0.1	Monitor

Metric	Value	Threshold	Action level
$\mathcal{D}_{\mathrm{KL}}$ **(amount)**	0.052	> 0.05	Investigate
KS test	$p_{\mathrm{val}} < 0.001$	$\alpha_{\mathrm{sig}} = 0.01$ (Bonferroni)	Investigate

Decision: Continue serving this fraud model with increased monitoring frequency (weekly instead of monthly), while investigating the transaction-amount shift. The moderate KL divergence and significant KS result justify that review under the team's configured triggers; the country PSI does not add an independent alert. These monitors establish only that the input distribution of transaction amounts has changed. The delayed labels must decide whether the fraud rule stayed fixed, a covariate shift that reweighting may address, or moved as well, a concept drift that requires adapting to recent outcomes. A model update becomes a candidate only if delayed labels show degradation beyond the team's validated tolerance and the investigation identifies a shift that retraining can address. The numerical cutoffs belong to this worked service; another model must calibrate its own thresholds from baseline variation, failure consequences, and response capacity.

14.4.2.5 Retraining decision framework

The fraud team kept its model in service despite a statistically significant shift because it separated a fired metric from a weight update. Generalizing that discipline across a fleet yields a three-level **retraining decision framework** that integrates drift statistics with performance monitoring so metric alerts do not become automatic retraining commands.

Level 1 automated monitoring The first level automates detection only. The fraud team's three triggers show its general form. Each threshold may start from a convention, such as the PSI bands, but is validated against the service's own historical baseline rather than copied unchanged. Any significance level is a per-test level after correction for the configured family of tests, so it must be recomputed when features are added to the family or its error budget changes. Crossing any threshold initiates an operational investigation rather than triggering an automated retraining pipeline.

Level 2 performance correlation Comparing scenario-calibrated alerts with observed performance yields three cases that guide investigation. If observed performance degradation exceeds the service's validated tolerance and coincides with detected drift, operators isolate the affected data slices to test whether recalibration, data repair, or retraining directly addresses the root cause. Conversely, when drift is detected while serving performance remains stable, the service continues operating under elevated monitoring while the team investigates the drift source. Finally, if performance degrades without detectable feature drift, operators investigate changes the monitored marginals could miss, including concept drift, joint-feature changes, and pipeline or model faults. An unchanged feature dashboard does not identify which of these caused the error.

Level 3 retraining vs. investigation Not all drift requires retraining. The evidence decides among three responses (a model update, investigation first, or continued monitoring), and each demands a different combination of statistical signal, observed performance, and external context (table 14.5).

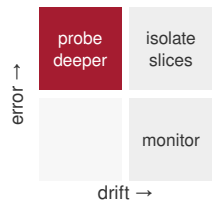

Degradation without drift is the case the feature dashboard cannot explain.

Table 14.5: Retraining Decision Rules: Drift response depends on the strength of the statistical signal, observed performance change, and external evidence about whether the environment has truly changed.

Action	Trigger conditions	Response logic
Update model	Performance exceeds a validated error tolerance, labels confirm a shift, root-cause analysis shows a model update addresses it, and offline plus shadow evaluation clears the release gate	The evidence identifies recalibration or retraining as the appropriate response rather than merely detecting a population change

Action	Trigger conditions	Response logic
Investigate first	A scenario-calibrated drift alert fires while performance remains stable, the shift is localized to nonpredictive features, or a seasonal or one-time event is plausible	The signal justifies diagnosis, but not a predetermined remediation
Continue monitoring	Drift statistics remain inside the service's calibrated baseline, performance stays within tolerance, and no external evidence suggests environmental change	The evidence does not yet justify changing the model, but the baseline remains under observation

The three-level framework turns drift detection from reactive troubleshooting into planned model maintenance as production environments evolve (Gama et al. 2014). No PSI band is a universal retraining rule: a large shift in a nonpredictive feature may be harmless, while a small shift in a safety-critical subgroup may demand immediate containment.

14.4.3 Monitoring and adaptation strategies

The fraud example compared two features of one service against one baseline. A fleet runs the same comparison for thousands of features across many replicas, regions, and model versions, so the design of the monitor itself determines what the fleet can detect and what detection costs. The decision rules above still apply; the new questions are which statistic to compute where, and what to change once the evidence warrants a model update.

Each statistic carries a computational and sampling price, and the price decides where and how often it runs. In this illustrative H100-class monitoring scenario, one MMD comparison with radial basis function kernels, which score two inputs as more similar the closer they lie ($\gamma = 1.0$), over a window of 10,000 samples takes an assumed 150 ms per window. Run hourly on sampled traffic, that is a small fraction of one accelerator's time, but the direct estimator's cost grows with the square of the window size, so doubling the window to gain detection power roughly quadruples the cost. The per-feature two-sample KS tests use 1,000 samples per window in this scenario. Each costs $\mathcal{O}(n \log n)$, dominated by sorting its window, which is cheap enough to run on every window for every feature. These tests, however, miss joint changes that preserve each feature's marginal distribution. The design therefore runs the cheap marginal tests continuously and reserves the joint kernel test for periodic, sampled checks. Neither sample count is a universal detection minimum; power depends on the shift and significance level. The service in this example uses PSI values from 0.1 to 0.25 as investigation triggers, not as universal statistical thresholds.

High-dimensional and temporal traffic requires more than one fleet-wide scalar. Marginal feature monitors can remain flat while correlations, sequence structure, or a small cohort changes. A production design therefore combines cheap per-feature alarms with embedding- or kernel-space tests, cohort slices, rolling windows at several durations, and change-point checks. Where these statistics run, on each replica or centrally from exported samples or sketches, sets both the telemetry load on the C^3 communication axis and the evidence the detector sees. The windows must preserve seasonality and event context: comparing Monday traffic with Sunday traffic can manufacture an alert, while averaging a week can hide an abrupt one-hour failure. Monitor design is part of the hypothesis, not a neutral wrapper around the model.

↩ The fleet monitoring hierarchy that decides where drift statistics run and which signals reach an owner is established in section 12.6.3.

When the Level 3 rules do call for a model update, the update itself is a budgeted response rather than an automatic retraining command. Online learning updates a model from sequential observations (Shalev-Shwartz 2012); preserving earlier behavior remains a separate design objective. The adaptation budget depends on model size, drift rate, and feedback latency: updating too aggressively can chase noise, while updating too slowly lets performance drift. In production, online-learning systems usually bound update frequency, state size, and serving latency explicitly rather than assuming adaptation is free. An update that fits the new distribution can also erase behavior the old one required, the catastrophic forgetting that also constrains on-device adaptation. Techniques like Elastic Weight Consolidation reduce it by penalizing changes to parameters important for previous tasks (Kirkpatrick et al. 2017).

↩ Catastrophic forgetting and the continual-adaptation techniques that limit it are established in section 11.8.2.

Adaptive ensemble methods maintain multiple models or hypotheses and weight or select among them using recent performance, making them useful under gradual concept drift (Gama et al.

2014). This approach trades extra serving and monitoring complexity for the ability to respond when a single static model no longer matches the deployment distribution. When the drifted data cannot leave the devices or regions that produce it, the same update can run as federated learning (section 11.6), which moves the adaptation to the data. The drift response then pays in communication rounds and transmitted bytes as well as training work, so the time from a confirmed shift to a validated update must be measured on the federated path rather than assumed from centralized retraining.

14.4.4 Per-request competence and selective prediction

The drift monitors and adaptation paths above ask whether traffic has changed in aggregate, yet a stable daily PSI can coexist with a rare, dangerous out-of-distribution request, while a large population shift can leave individual predictions well supported. Because a serving system must decide on every request whether that particular input lies within the model's demonstrated competence, fleet robustness needs a per-request decision contract alongside its population monitors.

That contract combines three distinct signals because no single score answers the whole question. Calibration tests whether a score of 0.8 corresponds to approximately 80 percent correctness on a defined validation population (Guo et al. 2017). Out-of-distribution detection scores how unlike the request is from the model's validated input distribution (Hendrycks and Gimpel 2017), because a calibrated classifier can remain confidently wrong on unfamiliar inputs. An OOD detector can flag novelty without estimating task error, so selective prediction turns both signals into an action policy, answering above the acceptance threshold, routing uncertain cases to a stronger model or reviewer, and abstaining when neither path can satisfy the risk bound (figure 14.6).

Let $\kappa(x)$ be an acceptance score, such as a calibrated confidence or a score that combines confidence with an OOD score, whose larger values indicate stronger evidence for answering, τ its threshold, $f(x)$ the prediction, and $\mathcal{L}(f(x), y)$ its loss against the observed outcome. For accepted set $\mathcal{A}_\tau = \{x : \kappa(x) \geq \tau\}$ at threshold τ, coverage is $\Pr(x \in \mathcal{A}_\tau)$ and selective risk is the mean loss conditioned on acceptance:

$$\text{coverage}(\tau) = \Pr(x \in \mathcal{A}_\tau), \qquad \text{risk}(\tau) = \mathbb{E}[\mathcal{L}(f(x), y) \mid x \in \mathcal{A}_\tau].$$

Sweeping τ produces a risk-coverage curve rather than one accuracy number. Raising the threshold reduces coverage, but lowers accepted-request risk only when the score orders requests by error risk well enough. That relationship must be measured on the deployment population, and the rejected traffic does not disappear. The fallback path needs enough queue, reviewer, or secondary-model capacity to absorb the abstention rate within the original latency and safety budget. If labels arrive days later, the system can monitor scores and fallback load immediately but cannot confirm selective risk until outcomes return. During that delay, rate limits and conservative thresholds can limit exposure, while the default response still requires validation against the application's hazards.

14.4.4.1 Finite-sample coverage with conformal prediction

Calibration describes whether predicted probabilities match empirical frequencies, whereas selective prediction chooses which requests the system will answer. Conformal prediction provides a different contract: it wraps a fitted predictor with a nonconformity score $s(x, y)$ and a held-out calibration set. Here the score measures how poorly a candidate outcome y agrees with input x, such as one minus the softmax probability assigned to class y, so larger values mean less conformity and make this score distinct from the acceptance score above.

On the n_{cal} held-out calibration examples, the true label y_i is known, so the system evaluates $s(x_i, y_i)$ once per example at its true outcome to measure how unusual true labels look under the fitted model. At serving time, the true label for a new input x is unknown, so the system trials every possible candidate label y in $s(x, y)$ and retains every class whose score does not exceed the cutoff calibrated on the held-out examples. For a skin-lesion classifier, the resulting prediction set might contain both benign and malignant when neither diagnosis looks sufficiently nonconforming, rather than force a single guess.

The target miscoverage rate is ϵ. For the finite threshold below, assume $1/(n_{\text{cal}}+1) \leq \epsilon < 1$. For n_{cal} calibration scores sorted in ascending order, set $k_\epsilon = \min\{\lceil (n_{\text{cal}}+1)(1-\epsilon) \rceil, n_{\text{cal}}\}$ and $\hat{q}_{1-\epsilon} = s_{(k_\epsilon)}$. The prediction set is

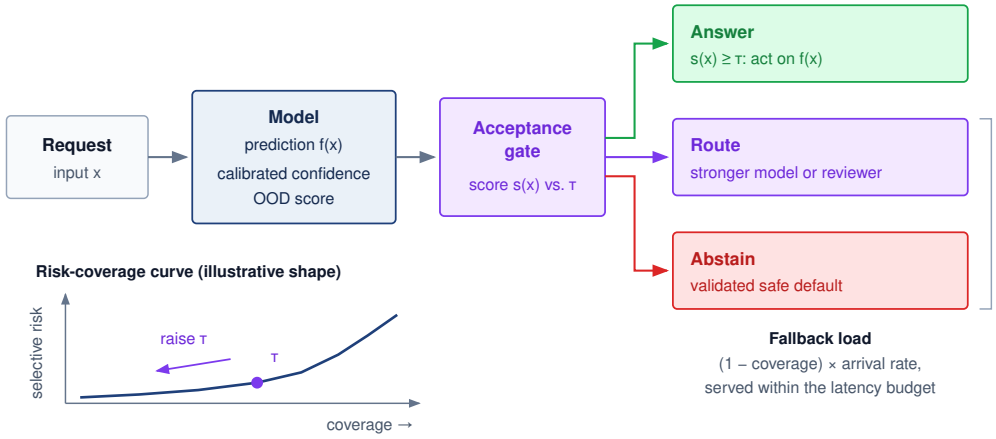

Figure 14.6: Per-Request Competence Routing: An acceptance gate compares each request's score with threshold τ and sends it to one of three paths. Accepted requests are answered from the model's prediction; the rest are routed to a stronger model or reviewer or receive a validated safe default. Raising τ moves the operating point down the risk-coverage curve, lowering selective risk only if the score orders errors well, and every point of coverage given up becomes fallback load that those paths must serve within the latency budget. The curve's shape is illustrative.

$$\mathcal{C}(x) = \{y : s(x,y) \le \hat{q}_{1-\epsilon}\}, \qquad \Pr(y_{n+1} \in \mathcal{C}(x_{n+1})) \ge 1-\epsilon.$$

Here the test pair (x_{n+1}, y_{n+1}) follows the $n = n_{\text{cal}}$ calibration examples. Exchangeability means that permuting their order does not change their joint distribution; i.i.d. sampling is sufficient. Under this assumption and the stated range of ϵ, the $(n_{\text{cal}}+1)$ rank correction makes this a finite-sample *marginal* coverage guarantee. Exchangeability makes the test point's score equally likely to take any of the $n_{\text{cal}}+1$ ranks among the combined scores. When averaging over the calibration sample and the new test example, the set therefore contains the true outcome at least at the target rate (Lei et al. 2018). Using the ordinary empirical $(1-\epsilon)$ quantile instead can undercover at finite sample sizes. The guarantee is not conditional on a particular realized calibration set, input, or subgroup. A rare subgroup can have poor coverage even when aggregate coverage is correct, particularly when it contributes few calibration examples.

The systems obligation is to preserve the conditions behind the guarantee. A distribution change between calibration and serving can invalidate the exchangeability argument; calibration age alone does not establish such a change. Even if observed coverage remains adequate, the original guarantee needs its assumptions rechecked. Repeatedly tuning the model, score, or threshold against the same calibration outcomes makes the held-out evidence adaptive rather than independent. Production systems therefore version the predictor, score, calibration snapshot, and target ϵ together; monitor realized coverage by consequential slice once labels arrive; and refresh or quarantine the contract when traffic changes. Conformalization does not protect against arbitrary shift.

Set size or interval width supplies an operational load signal. Larger outputs convey less specific predictions and can trigger abstention, a stronger model, or human review, so their tail distribution determines fallback capacity just as the risk-coverage curve determines rejected traffic. The target coverage is therefore not free: increasing it usually widens sets, raises review demand, or reduces the fraction of requests on which the system can act automatically.

The failure modes above arise without an adversary, but the boundary with deliberate manipulation is porous. Distribution monitors can themselves be evaded by inputs crafted to stay under their thresholds, turning a defensive tool into a blind spot. Once an adversary stops relying on natural variation and begins optimizing against the model's decision boundaries to force specific errors, the threat model moves from environmental robustness to input-level attacks.

Checkpoint 14.2: Drift detection and the retraining decision

Environmental shifts split into distinct types, and the chapter's drift framework deliberately separates a fired metric from a retraining command.

Distinguishing the shift type

- ☐ **Shift taxonomy**: distinguish covariate shift, label shift, and concept drift by which probability the change touches, and explain why feature-statistics monitoring stays silent under pure concept drift.
- ☐ **Shortcut failure**: identify the failure in a cow-on-a-beach misclassification where the world has not changed but the model learned the wrong lesson, and explain why standard empirical risk minimization encourages it.

Reasoning about thresholds and cost

- ☐ **Threshold judgment**: select retraining, investigation, or continued monitoring for a high-importance feature whose population stability index is in the major-shift band while precision and recall stay flat.
- ☐ **Cost asymmetry**: explain why concept drift can cost more to remediate than covariate shift when adapting the conditional relationship requires fresh labels.

Deciding request by request

- ☐ **Per-request contracts**: explain when raising a selective-prediction threshold lowers the risk of accepted requests, and what conformal coverage at the target rate does and does not promise for a rare subgroup.

14.5 Adversarial Attacks

For a classifier that labels skin-lesion images, an attacker who modifies an image of a benign lesion can induce a malignant prediction while preserving its apparent clinical content. Whether that manipulation succeeds depends on the classifier, the allowed modifications, and the evaluation conditions. The high-dimensional decision boundaries learned by deep neural networks possess counterintuitive blind spots, and adversarial attacks expose them by testing what an adversary can force when controlling gradients, queries, physical sensors, or training data.

Definition 14.3: Adversarial attack

Adversarial attack is a deliberate manipulation of a model's input, interaction, or environment designed to violate a specified behavioral objective. The manipulation may be norm-bounded and visually subtle, but it can also be semantic, physical, query-based, or conspicuous.

1. **Significance**: Adversarial evaluation asks whether an attacker can choose inputs or transformations that drive the model outside its intended behavioral envelope. High-dimensional, norm-bounded image perturbations are one important case, but patches, rotations, printed objects, malicious text, and repeated black-box queries expose different system boundaries.
2. **Distinction**: Unlike random corruption sampled independently of the target, an adversarial manipulation is selected to cause failure. It need not use gradients: attackers may use surrogate models, decision-only queries, search, semantic edits, or control over the physical sensing environment.
3. **Common pitfall**: A frequent misconception is that adversarial vulnerability is a "bug" to be patched. In reality, it is a structural vulnerability of many standard neural networks trained by empirical risk minimization; robust defense typically requires fundamental changes to the objective function (for example, adversarial training).

When an attacker can inspect a model's gradients directly, a change no person can see is enough to fool an image classifier. A perturbation that follows the sign of the loss gradient ($\epsilon_{\text{adv}} = 0.007$), using the fast gradient sign method defined in section 14.5.1, flips GoogLeNet's prediction (figure 14.7) from panda (57.7 percent confidence) to gibbon (99.3 percent confidence). The classifier is left paradoxically more confident in the erroneous label than in the ground truth.

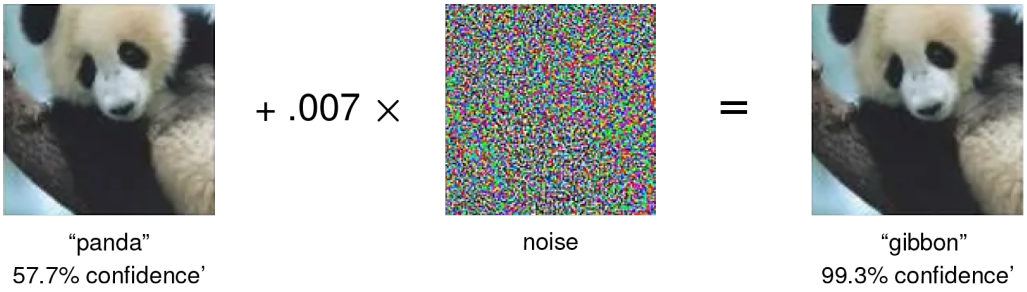

Figure 14.7: Imperceptible Misclassification in GoogLeNet: Adding a perturbation of $\epsilon_{\text{adv}} = 0.007$ in the sign direction of the loss gradient changes GoogLeNet's prediction from panda (57.7 percent confidence) to gibbon (99.3 percent confidence), although the two images look identical to a person. The attack does not merely flip the label; it leaves the model more confident in the wrong answer than it was in the right one. Source: (Goodfellow et al. 2015).

A visually small change need not produce a small change in a classifier's score.[7] Adversarial perturbations exploit how input changes accumulate in the model's decision function. Even a linear function can be sensitive to many small, aligned changes across input dimensions, so nonlinearity alone does not explain the vulnerability. For a linear score, changing every input coordinate by the full per-coordinate budget in the direction of its weight's sign moves the score by that budget times the sum of the absolute weights. For weights of typical size, that total grows in proportion to the number of input dimensions, even though no single coordinate moves visibly. The relevant constraint is whether the allowed perturbation can move the input across the learned decision boundary, and how an attacker finds that direction depends on the access the deployed system exposes.

[7] **Human vs. machine perception**: First highlighted by Szegedy et al. (2013), neural networks learn statistical correlations in pixel space rather than the semantic invariances human vision enforces. Agreement on clean examples does not establish agreement under every small image change. Evaluation must therefore specify which transformations should preserve the label and test whether the model preserves that behavior; gradient-based training alone neither explains every attack nor guarantees the required invariance.

14.5.1 Attacks by attacker access

An attacker who can read a model's weights and gradients can flip a prediction in a single backward pass, whereas an attacker facing a rate-limited API must spend thousands of queries or train a local surrogate model to find a boundary-crossing input. The useful axis is therefore not the attack name by itself, but the attacker's access. Each mechanism reveals a different cost that the defense must raise: gradient access, optimization time, surrogate-model construction, or physical control over the sensor environment.

The most direct case is white-box gradient access. Neural networks compute gradients to learn how parameter changes reduce loss; an attacker with access to gradients can run that logic against the input instead. In the GoogLeNet panda example of figure 14.7, the gradient with respect to the input image gives a local, first-order estimate of how pixel changes affect the loss. The **fast gradient sign method (FGSM)** turns that idea into a single-step attack by choosing each feature's perturbation sign to maximize that first-order increase within a coordinate-wise budget.[8]

FGSM computes the adversarial example in one step:

$$x_{\text{adv}} = x + \epsilon_{\text{adv}} \cdot \text{sign}\big(\nabla_x \mathcal{L}(\theta, x, y)\big)$$

In this formulation, x denotes the original input, x_{adv} is the generated adversarial example, and $\mathcal{L}(\theta, x, y)$ represents the prediction loss evaluated at model parameters θ and true label y. The gradient term $\nabla_x \mathcal{L}$ describes local loss sensitivity, $\text{sign}(\cdot)$ extracts each component's sign to maximize the linearized loss within the coordinate-wise budget, and the scalar ϵ_{adv} bounds the maximum allowable perturbation magnitude under the ℓ_∞ norm. This single step need not maximize the nonlinear loss or cause misclassification. Image implementations also clip the perturbed input to

[8] **Fast gradient sign method (FGSM)**: Proposed by Goodfellow et al. (2015) at ICLR, FGSM generates adversarial examples in a single gradient step, making it practical to use during training. This dual role is its lasting systems significance: the same attack that exposed neural network fragility became a simple adversarial-training primitive, where FGSM-generated perturbations augment clean batches with examples chosen to increase loss.

its valid pixel range. The subscript distinguishes this adversarial radius from other uses of ϵ in the book, such as the differential-privacy budget.

Geometrically, moving along the sign vector of $\nabla_x \mathcal{L}(\theta, x, y)$ carries input x to a corner of the bounding cube defined by the ℓ_∞ budget ϵ_{adv}, producing x_{adv}. In the illustration of figure 14.8, that corner lies across the decision boundary in the adjacent class region. A different boundary or gradient need not produce a successful attack within the same budget.

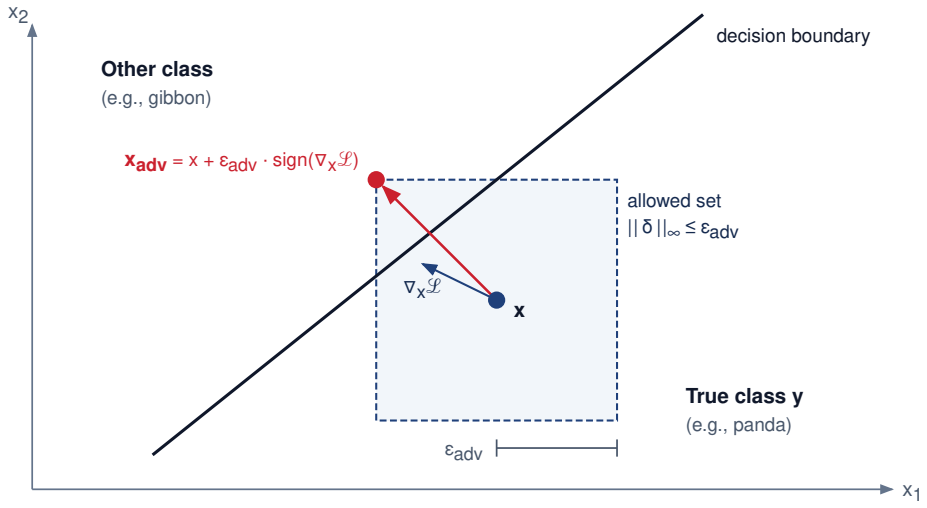

Figure 14.8: Gradient-Based Attack Geometry: Around a correctly classified input x, the dashed square marks every perturbation allowed by the ℓ_∞ budget ϵ_{adv}. FGSM moves each coordinate by the full budget in the sign direction of the loss gradient $\nabla_x \mathcal{L}(\theta, x, y)$, which lands x_{adv} on a corner of that square, the allowed point that maximizes the first-order increase in loss. A step no larger than ϵ_{adv} in any coordinate therefore carries the input across the decision boundary. The class labels are illustrative.

FGSM is cheap because it takes one step. The **projected gradient descent (PGD) attack** (Madry et al. 2018) spends more compute for a stronger attack: it repeatedly applies gradient updates and projects each step back into the allowed norm ball around the original input. That iterative refinement makes PGD a standard white-box robustness benchmark. The **Jacobian-based saliency map attack (JSMA)** (N. Papernot et al. 2016) uses the Jacobian to identify the most influential input features and perturb a smaller set of dimensions toward a target class. These methods are most effective in white-box settings,[9] where the attacker knows the model architecture and gradients.

When the attacker cares less about speed and more about stealth, the attack becomes an optimization problem. The Carlini and Wagner (C&W) attack[10] (Carlini and Wagner 2017) searches for a small perturbation that causes misclassification under a chosen distance measure. Minimizing that distance encourages similarity, but does not certify human perceptual equivalence or a globally minimal perturbation. Instead of merely following the sign of a gradient, C&W optimizes a custom objective that trades perturbation size against confidence in the wrong answer.

The attack can search under different norm constraints, such as ℓ_2 or ℓ_∞. Its perturbations can be visually subtle, although a perturbation a person cannot see may still be one a detector flags. The **elastic net attack to DNNs (EAD)** adds elastic net regularization, combining ℓ_1 and ℓ_2 penalties to generate sparse, localized perturbations. **DeepFool** (Moosavi-Dezfooli et al. 2015) takes a geometric route to the same goal: it linearizes the classifier around the input, steps to the nearest boundary of that linear approximation, and repeats until the label flips, giving a fast estimate of the smallest perturbation that suffices. These optimization-based methods are more computationally intensive than gradient-sign attacks, but they give the attacker finer control over the adversarial example's geometry.

Black-box attackers lose direct gradient access but keep two routes. Transfer attacks rely on transferability, the tendency of an adversarial example to fool models other than the one it was

9 **White-box attack**: Adversarial attack with complete model knowledge (architecture, weights, gradients); methods like PGD and Carlini-Wagner (C&W) are strong because they optimize directly against the model rather than probing it from the outside. Though less realistic than black-box scenarios for many deployed systems, white-box analysis establishes a demanding benchmark within a specified norm, perturbation budget, and attack procedure. Passing PGD is evidence for that threat model, not a universal guarantee against every weaker-looking attack.

10 **Carlini and Wagner (C&W) attack**: Proposed in 2016 (IEEE S&P 2017), C&W formulates adversarial example generation as a constrained optimization problem that searches for a small perturbation causing misclassification without guaranteeing a global minimum. Its significance is methodological: C&W broke defensive distillation (an early defense that masked gradients to blunt weaker attacks) and other evaluated defenses, establishing the principle that robustness claims must be evaluated against optimization-based attacks, not merely gradient-sign heuristics.

crafted on. The attacker crafts examples offline on a surrogate model and submits them to the target, typically through its public API, without seeing its weights or gradients. Query-based attacks probe the target itself. A decision-based, or boundary, attack needs only the returned label. It starts from an input the model already misclassifies and walks along the decision boundary toward the original input, querying at each step to stay on the wrong side. Its cost is the number of queries, which is why query budgets and rate limits belong to the defense.

↩ The adversary knowledge spectrum, from white-box gradient access to black-box attackers who work through a surrogate or a public interface, is established in section 13.4.1.2.

Transfer works because models trained for the same task on similar data tend to learn similar decision boundaries, so a perturbation that pushes an input across one model's boundary often pushes it across another's. Transfer success depends on model similarity, training-data overlap, and regularization, and attackers can improve it with input diversity, such as random resizing or cropping, and with momentum during optimization. Closing that route costs the defender training work. Ensemble adversarial training trains against perturbations generated on several models to improve transfer robustness, but the extra attack generation and model diversity make it a deliberate training-budget decision rather than a free mitigation (Tramèr et al. 2017).

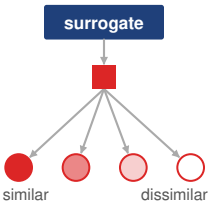

An example crafted on a surrogate fools similar models more often than dissimilar ones.

14.5.2 Physical-world attacks

White-box and black-box digital attacks modify pixel tensors directly in memory, whereas an attacker in the field cannot edit the camera buffer and must place the perturbation on physical objects such as road signs, clothing, or 3D-printed surfaces. **Adversarial patches** are printed patterns placed in the scene so that cameras and detectors misread it even after the pattern passes through sensors, distance, lighting, viewing angle, and ordinary deployment variation. The modification need not be visually imperceptible, because graffiti-style markings and sticker patches that pass for ordinary vandalism can still induce misclassification through the camera pipeline (Eykholt et al. 2017). The robustness question therefore becomes operational rather than purely mathematical, because a defense that works on a saved image may fail once the attack passes through a camera lens, compression, motion blur, and changing illumination. These threats matter most for AI systems deployed in physical spaces, such as autonomous vehicles, drones, and surveillance systems, where a model error becomes a safety, security, and accountability problem in the world.

> Example 14.1: The stop sign attack
>
> **Scenario**: The road-sign sticker attack places black and white stickers on a stop sign so that a vision model reads it as a speed limit sign (figure 14.9) (Eykholt et al. 2018). The robustness question is what its two reported success rates say about how an evaluation must be run.
>
> **Diagnosis**: The attack succeeded on every tested laboratory image but on 84.8 percent of frames in the moving-vehicle field test. Each number is evidence only for its own capture conditions, and the gap between them does not isolate any one factor, such as distance, viewing angle, lighting, or motion.
>
> **Systems lesson**: A physical attack must be evaluated through the full sensor loop, not on saved benchmark images. The evaluation should optimize the attack averaged over random draws from the expected range of transformations, a method called expectation over transformation, and then test it with real captures. Any claimed mitigation, such as checking consistency across frames (spatio-temporal validation) or across sensors (sensor fusion), must be tested against attacks on the complete perception stack.

↩ The road-sign sticker attack is established as a security threat in section 13.5.

The attacker's access determines the evaluation a defense must survive. Gradient access, optimization, surrogate transfer, and physical sensor manipulation each call for their own evaluation (table 14.6), so a result under one attack mechanism does not become a claim about all of them.

Figure 14.9: Adversarial Stickers on Stop Signs: In the cited study, the modified signs achieved targeted misclassification on all tested laboratory images and 84.8 percent of moving-vehicle field-test frames. The gap between those conditions shows why physical attacks must be evaluated across the actual sensing path. Source: (Eykholt et al. 2018).

Table 14.6: Adversarial Attack Categories: Attack mechanisms determine defense selection: gradient and optimization attacks require robustness against explicit decision-boundary crossing, transfer attacks test black-box exposure, and physical attacks require evaluation through the deployed sensor loop.

Category	Method	Mechanism
Gradient	FGSM	Perturbs inputs along the loss gradient
	PGD	Iterative multi-step FGSM refinement
	JSMA	Targets the most influential features
Optimization	C&W	Minimizes perturbation size subject to misclassification
	DeepFool	Estimates a small boundary-crossing perturbation by local linearization
	EAD	Elastic net regularization for sparse perturbations
Transfer	Transferability	Adversarial examples transfer across models (black-box)
Physical	Patches	Printed patches fool detectors in the real world
	3D Objects	Sculpted objects deceive sensors in deployment

Medical-imaging attacks can cause misdiagnosis (Tsai et al. 2023), while fraud and trading models face analogous manipulation risks. The common vulnerability is reliance on features that remain predictive on clean data but fail under crafted perturbations (Goodfellow et al. 2015; Madry et al. 2018). Defense shifts the resource cost: adversarial training adds training work (Bai et al. 2021), while runtime detection such as feature squeezing adds inference evaluations (Xu et al. 2018). The choice among these spending points motivates the defenses in section 14.6.

14.6 Adversarial Defenses

The classifier that called a panda a gibbon in figure 14.7 was not broken by random noise; a perturbation budget of 0.007 per pixel, spent along the gradient sign, was enough. Defending a deployed perception model against that attacker requires a threat model that names the attacker's access and a budget for testing and mitigation. The resulting robustness claim must state what the system pays in clean accuracy, training compute, inference latency, or operational coverage.

14.6.1 Adversarial defense workflow

A team that applies adversarial training without first naming the perturbation norm and budget can multiply training compute several-fold and lose clean accuracy while leaving the actual deployment threat uncovered. Preventing that misallocation requires an adversarial-defense workflow with four ordered layers:

1. **Define the threat model**: Specify the perturbation budget, attacker knowledge, and whether the attack occurs at inference time or during training.
2. **Choose where to spend the defense budget**: Decide whether to spend it in training, certification, input detection, or serving-time guardrails.
3. **Measure the trade-off**: Evaluate both sides of the defense, because improving worst-case behavior can reduce clean accuracy, raise inference latency, or multiply training cost.
4. **Keep evaluation adversarial**: Test the defense against attacks stronger than the ones used to design it.

The order matters because spending and evaluation choices made before the threat model is specified cannot show protection against the attacker the deployment actually faces. For model-facing perturbation threats, the most direct training-time defense is **adversarial training**, which incorporates adversarial examples into the training process itself. The mechanism reuses the attacks of the previous section as a data generator. Each training step augments the batch with perturbed copies of its own inputs, so the model is trained on the high-loss points an attacker would choose. In the simplest form, FGSM generates the perturbations and each batch mixes clean and adversarial examples (listing 14.1).

Listing 14.1: Adversarial Training Pseudocode: Framework-neutral adversarial training using FGSM-style perturbations during training, mixing clean and perturbed data to improve robustness against gradient-based attacks.

```
function adversarial_training_step(model, clean_batch, labels, epsilon):
    logits = model(clean_batch)
    clean_loss = loss(logits, labels)

    input_gradient = gradient(clean_loss, clean_batch)
    perturbation = epsilon * sign(input_gradient)
    adversarial_batch = stop_gradient(
        clip(clean_batch + perturbation, valid_input_range))

    mixed_batch = concatenate(clean_batch, adversarial_batch)
    mixed_labels = concatenate(labels, labels)

    return loss(model(mixed_batch), mixed_labels)
```

In listing 14.1, the sign step bounds each coordinate's change by the perturbation budget, while clipping keeps the perturbed input inside the valid input range (for example, $[0, 1]$ for normalized pixels). Concatenation doubles the effective batch size. The stop-gradient operation treats generated adversarial inputs as fixed during the outer weight update, rather than differentiating through attack construction. The caller differentiates the returned loss with respect to the model parameters and performs the optimizer update. The perturbation budget ϵ_{adv} (the `epsilon` argument) bounds the allowed input change, while the attack algorithm determines how effectively it searches that set; computation reuse can reduce its training overhead (Shafahi et al. 2019).

Adversarial example generation increases training time because each training step requires additional forward and backward passes through the model (Madry et al. 2018; Bai et al. 2021). Memory overhead depends on whether the implementation stores clean and adversarial examples together, recomputes perturbations, or reuses gradient information. Iterative attacks like PGD add multiple optimization steps, so the implementation must budget that gradient work; optimized variants reduce overhead by reusing computations rather than treating every attack step as a separate full training pass (Shafahi et al. 2019). A robustness recipe that increases model capacity also adds memory and training cost. The perturbation budget and training recipe require matched clean

and adversarial validation, and the perspective below shows why that matching decides what the robustness penalty actually is.

Systems Perspective 14.1: Pricing the robustness penalty correctly

Suppose a standard ResNet-50 records 76 percent Top-1 accuracy on clean ImageNet, while an adversarially trained model records about 50 percent under a specified ℓ_∞ attack at $\epsilon_{\text{adv}} = 4/255$, the budget conventionally used for ImageNet. Subtracting these values yields 26 percentage points, but that gap compares different models under different evaluation conditions. It is not the clean-accuracy cost of robustness.

The clean-accuracy penalty must compare standard and robust models on the same clean test set. The robustness benefit must compare those models under the same attack, perturbation budget, and attack strength. The four cells are the standard model on clean data, the standard model under attack, the robust model on clean data, and the robust model under attack. The penalty is the first minus the third, the benefit is the fourth minus the second, and the quoted pair supplies only the first and the fourth. Reporting all four cells prevents an apparent trade-off from mixing a model effect with an evaluation-condition effect.

Systems insight: Robustness can reduce clean accuracy and increase compute, but each cost needs a like-for-like baseline. A deployment gate should report clean accuracy, empirical robust accuracy, attack specification, and evaluation cost separately.

The robustness compute penalty (principle 24) quantifies this cost: PGD-style adversarial robustness can demand several times more compute per epoch than standard optimization because each batch runs inner attack steps. In Fleet Law terms (equation 1.3) the multiplier lands on the per-device compute term T_{compute}. At a fixed worker count it adds local work, but the elapsed-time penalty depends on reuse and on which term limits the step. Adding workers changes communication and synchronization costs with N and may not recover the original duration. Notebook 14.2 makes the multiplier concrete.

Napkin Math 14.2: The cost of defense

Problem: A team is training a robust classifier for an autonomous vehicle. Instead of the single FGSM step in listing 14.1, it uses adversarial training with a seven-step PGD attack (PGD-7) that searches for a high-loss perturbation inside a specified norm ball for every training batch. Following the standard PGD adversarial-training recipe (Madry et al. 2018), it trains on the adversarial batch alone rather than on the listing's doubled clean-plus-adversarial batch. How much does this extra attack-generation work slow down the training run?

Math: In this example, generating the adversarial input uses K_{PGD} additional gradient steps per training sample, each a forward and backward pass with respect to the input, before the one weight-update pass on the adversarial batch. The estimate treats each attack step as equal in cost to the standard training pass, with no reuse; it is a work model rather than a measured wall-clock slowdown. By the same accounting, the FGSM listing spends one attack pass plus an update pass over a batch twice the original size, about three standard passes in all, so PGD-7 trades a further factor of more than two in training work for a stronger inner attack.

1. **Forward/backward passes**: 1 (Standard) + 7 (Attack Generation) = 8 total passes.
2. **Training slowdown**: 8× under the equal-cost-pass assumption.
3. **Utility evidence**: The added work must be justified by measured robust accuracy under the specified attack. Generalization to other attacks requires separate evaluation. Pricing any clean-accuracy cost requires the like-for-like comparison in perspective 14.1, not a subtraction of robust accuracy from another model's clean accuracy.

Systems insight: In this PGD-7 work model, the team budgets 8× the baseline compute before measuring the robustness benefit or actual elapsed-time penalty. In a machine learning fleet,

> robustness is not a setting that can be switched on; it is a budget the team spends, and the budget must be justified by the failure consequence of the component being hardened.

For lower-risk components, external guardrails such as input filtering and output verification can cost less than training intrinsic robustness into the weights, provided adaptive evaluation (section 14.6.1.3) supports acceptable residual risk against an attacker who knows the guardrail. Passing a finite set of attacks does not prove that no bypass exists. Safety-critical paths more often justify the training and certification budget.

14.6.1.1 Certified defenses

Adversarial training provides empirical resistance that must be tested against the deployment threat model. It can generalize beyond the examples generated during training, but attack evaluation alone does not prove invariance throughout an allowed perturbation set. **Certified robustness** offers a mathematical guarantee: for a given input x and radius ϵ_{adv}, no perturbation $\|\delta_{\text{adv}}\|_p < \epsilon_{\text{adv}}$ exists that changes the model's prediction. This certifies local prediction stability, not that the unchanged prediction is correct. The guarantee formalizes the intuition that a robust model limits how quickly its outputs can change as its inputs change. The model's **Lipschitz constant**[11] bounds that rate. If no class score can move by more than the Lipschitz constant times the perturbation size, the margin between the top score and the runner-up score can shrink by at most twice that amount, so the predicted label holds for any perturbation smaller than the margin divided by twice the Lipschitz constant. The certified radius is a provable lower bound on the input-perturbation part of the robustness margin that compression and other efficiency changes must preserve.

[11] **Lipschitz continuity**: A mathematical property that bounds how much a function's output changes relative to its input change ($\|f(x) - f(x')\| \leq K\|x - x'\|$). A small Lipschitz constant (K) limits score movement but does not by itself prevent a label flip. Reducing K can enlarge the certified radius, though enforcing a low constant can trade away clean-data accuracy and add training cost.

Computing or constraining a tight Lipschitz constant directly across a deep network is difficult because layer-by-layer norm bounds multiply across depth and quickly become too loose to certify useful radii. A widely used technique that avoids layerwise bound propagation on high-dimensional inputs like ImageNet is **randomized smoothing**. Instead of classifying x directly, it classifies the smoothed function $g(x)$, defined as the most probable class predicted by the base classifier f under Gaussian noise: $g(x) = \arg\max_j \Pr(f(x + \delta) = j)$ where $\delta \sim \mathcal{N}(\mathbf{0}, \sigma^2\mathbf{I})$. Because two Gaussian distributions centered at x and $x + \delta_{\text{adv}}$ overlap heavily whenever $\|\delta_{\text{adv}}\|_2$ is small relative to the noise scale σ, averaging the base classifier over Gaussian draws analytically bounds how fast class probabilities can shift with δ_{adv}, regardless of how sharp the underlying neural network is internally.

Cohen et al. (2019) proved a tight multiclass bound for the certified radius R using a lower bound on the top-class probability p_A and an upper bound on the runner-up probability p_B: $R = \frac{\sigma}{2}\left(\Phi^{-1}(p_A) - \Phi^{-1}(p_B)\right)$, where Φ^{-1} is the inverse standard normal CDF. In the binary special case, where $p_B = 1 - p_A$, this reduces to $R = \sigma\Phi^{-1}(p_A)$. This Gaussian-smoothing bound certifies a Euclidean (ℓ_2) perturbation radius, measured in the same units as the input to which the noise is added; it does not automatically certify the other norms used in attack benchmarks. If a binary decision has top-class probability $p_A = 0.999$ under noise $\sigma = 0.5$, the simplified radius is approximately 1.5 in those input units. For an image scaled to $[0, 1]$, that radius bounds the Euclidean length of the whole perturbation vector. It is not directly comparable to the per-coordinate ℓ_∞ budget of the robustness-penalty comparison. The radius and that budget differ because an ℓ_∞ perturbation of size ϵ_{adv} spread over d pixels can have Euclidean length up to $\epsilon_{\text{adv}}\sqrt{d}$, so a per-pixel budget that looks tiny can correspond to a long Euclidean vector on a high-dimensional image.

The multiclass case rarely estimates the runner-up probability separately. Because the class probabilities sum to one, no competing class can exceed $1 - p_A$, so a lower bound on p_A alone also bounds every competitor. Estimating that lower bound from samples turns certification into a statistical estimation problem and adds a confidence level to the certificate; an observed class frequency alone is not a certified probability bound.

Certification also carries accuracy and compute trade-offs. On ImageNet, the same work reports nontrivial certified accuracy for randomized smoothing, including certification at radius 0.5, but certification requires sampling many noise vectors per inference to estimate class probabilities with sufficient confidence. The increase in inference latency makes the sample budget a deployment constraint. Two designs keep certification outside the tight serving budget. Asynchronous auditing certifies a sampled fraction of served requests offline, so the fleet learns how often its answers were

certifiable without delaying any of them. A selective safety interlock certifies only the requests routed to a safety-critical action, such as a braking decision, and lets the remaining traffic take the uncertified path.

Certified defenses cover one part of the adversarial workflow. Production systems still need detection paths for suspicious inputs, mitigation strategies that change training or serving behavior, and evaluation procedures that verify which threat models remain covered.

14.6.1.2 Detection techniques

Detecting adversarial examples before they reach the model forms the first line of defense only when the signal reaches a serving decision: reject, transform, route, or audit. A population-level signal asks whether a window of current inputs still resembles the reference population. The two-sample tests of section 14.4.2.3 serve here unchanged. A KS test, or the Anderson-Darling test, a variant that weights differences in the distribution tails more heavily, can flag a traffic window whose inputs no longer resemble the reference. Such an alert does not identify an individual request as adversarial; per-request rejection requires an additional decision rule. These tests also inherit the drift monitors' blind spot, since an adversarial perturbation can preserve each feature's marginal distribution while corrupting the joint structure that joint-distribution or embedding-space comparisons would expose.

An attack that preserves marginal distributions can evade that signal. A complementary per-request signal asks whether a prediction survives a simplifying transformation of its input. **Feature squeezing**[12] (Xu et al. 2018) reduces input-space complexity through color bit-depth reduction or spatial smoothing, then compares the model's prediction before and after the transformation. A large prediction change flags sensitivity to the transformation. It does not by itself establish malicious intent or identify which image detail caused the change; legitimate inputs can also trigger the detector.

Another signal asks whether the model is uncertain about its own answer. An estimate of model uncertainty[13] can route an input for rejection, human review, or a more robust model. Bayesian neural networks estimate uncertainty by placing a probability distribution over each weight and averaging predictions over sampled weights, so disagreement among those samples signals uncertainty. Ensemble methods compare predictions from independently trained models and use disagreement as the warning signal (Lakshminarayanan et al. 2017). Their detection value depends on whether uncertainty distinguishes harmful inputs from legitimate ones under the chosen attack. Repeated predictions add inference work, so deployments must evaluate both missed attacks and serving cost rather than assume that a more expensive estimator detects more.

Dropout, the training-time regularizer that randomly deactivates neurons (Hinton et al. 2012), offers a cheaper approximation to that weight sampling. It can be repurposed for uncertainty estimation through **Monte Carlo dropout**[14] at inference time, where multiple forward passes with different dropout masks approximate the uncertainty distribution. Monte Carlo dropout is itself an approximate Bayesian method. Its uncertainty estimates must be validated on the deployment population rather than ranked by the original purpose of dropout.

Beyond detection, adversarial training, and certification, the remaining mitigation strategies move the spending point rather than eliminating the cost. **Defensive distillation** (Nicolas Papernot, McDaniel, Wu, et al. 2016) teaches a student model from the teacher's soft labels. Its apparent adversarial protection came from gradient masking, in which a defense leaves adversarial inputs in place but makes the model's gradients too flat or uninformative for a gradient-based attack to find them. Attacks that bypassed the masked gradients defeated it (Carlini and Wagner 2017), so making an attack optimizer fail is not evidence that nearby adversarial inputs are absent. Input preprocessing spends the defense budget at serving time: image denoising, JPEG compression, random resizing, padding, and random transformations try to erase perturbations before the model sees them. Ensembles spend it through redundancy, combining models with different architectures, training data, or preprocessing paths so that a perturbation that fools one member is less likely to fool all of them.

14.6.1.3 Evaluation and testing

Evaluation closes the budget loop: the system must measure which attacks the defense actually covers and what performance it sacrifices under controlled attack conditions. Robustness metrics

12 **Feature squeezing**: In the original evaluation (Xu et al. 2018), squeezing 256 color levels to 16 or applying median filtering eliminated many adversarial examples while maintaining high clean accuracy. The same transformations can remove useful clean-input features and do not cover every attack mechanism, so adaptive attacks that anticipate the squeezer require separate validation before the defense is treated as production-ready.

13 **Model uncertainty (epistemic uncertainty)**: Uncertainty associated with limited knowledge of the predictive model, distinguished from aleatoric uncertainty in the observations. An estimate can support deferral, but it must be evaluated against actual errors and unfamiliar inputs; a low estimate does not certify competence. Repeated-pass or ensemble estimators add inference work, and any human fallback must meet the application's response deadline.

14 **Monte Carlo dropout**: Proposed by Gal and Ghahramani (2016), MC dropout reinterprets dropout as approximate Bayesian inference by keeping dropout active across repeated stochastic forward passes. The variance across predictions provides an uncertainty estimate with no architectural changes. Each additional sample adds inference work, so deployments must select the sample count against their latency and uncertainty budgets.

quantify resilience through accuracy on adversarial examples, the distortion required to fool the model, and performance over attack budgets. Standardized benchmarks such as MNIST-C (Mu and Gilmer 2019), CIFAR-10-C, and ImageNet-C (Hendrycks and Dietterich 2019) measure common corruptions rather than malicious worst-case behavior, so their results occupy a different cell of table 14.1.

Empirical attack evaluation must be adaptive to the defense. The evaluator should know the preprocessing, detector, randomness, ensemble, and rejection rule. It should optimize through differentiable components and, for randomized components, use expectation over transformation, here averaging the attack gradient over the defense's random draws rather than over physical transformations. It should test transfer and query-based attacks when gradients are unavailable, and include the complete pipeline rather than a saved model alone. Multiple random restarts and attack budgets can strengthen evidence within that threat model, but a passed PGD test proves neither that PGD found the inner maximum nor that semantic, physical, or differently normed attacks are covered. Only a certificate establishes a guarantee, and only for its stated input, norm, radius, and assumptions.

RobustBench approximates that discipline at scale with one standardized, parameter-free attack ensemble, AutoAttack, applied to every published defense under one fixed threat model, which for CIFAR-10 is a perturbation of at most 8/255 per pixel in the ℓ_∞ norm (Croce et al. 2020; Croce and Hein 2020). Across the 91 dated entries (figure 14.10), the best accuracy under attack rose from 44 percent in 2018 to 75 percent in 2024, while the same model kept 94 percent accuracy on clean images. Even the strongest defense therefore still fails on about 25 percent of attacked inputs. The four entries at the bottom of the plot are published defenses that keep high clean accuracy yet fall below 5 percent accuracy under the standardized attack, the result an evaluation that is not adaptive to the defense can miss.

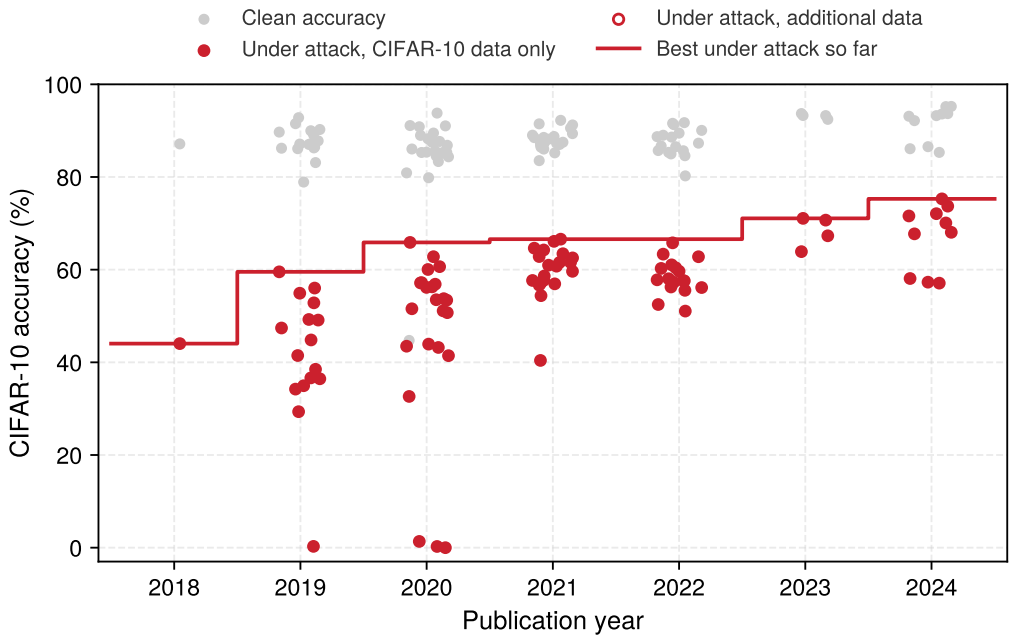

Figure 14.10: Standardized Robustness Evaluation of Published Defenses: Clean accuracy (gray) and accuracy under AutoAttack (red) for CIFAR-10 defenses on the RobustBench leaderboard by publication year, at a perturbation budget of 8/255 in the ℓ_∞ norm. Filled red points trained only on CIFAR-10; hollow points used additional data. The line tracks the best robust accuracy published by each year. Points near zero are defenses that do not withstand the standardized attack. Data from Croce et al. (2020).

Adversarial training, certified defenses, and input detection address attacks on model inputs at inference time within bounded threat models. They do not establish that the model was trained on

trustworthy data. If an adversary compromises the data pipeline long before the model is compiled, inference-time behavioral tests may reveal a backdoor, a trigger pattern planted through training data that makes the model output an attacker's chosen label, but cannot by themselves establish the integrity of the training pipeline. Poisoning attacks therefore demand their own class of defenses.

Checkpoint 14.3: Choosing and budgeting an adversarial defense

An adversarial defense is a budgeted response to a specific threat model. Stronger guarantees commonly change clean accuracy, compute, latency, or coverage, but the direction and magnitude depend on the workload, model, defense, and evaluation boundary.

Matching defense to threat

- ☐ **Worst-case evaluation**: explain why passing a white-box PGD test provides empirical evidence only for the evaluated model, norm, budget, restarts, and attack implementation rather than a general worst-case guarantee.
- ☐ **Threat-model fit**: match the defense emphasis to an API-facing classifier and a safety-critical road-sign detector, and explain why adversarial training alone does not transfer between them.

Reasoning about the budget

- ☐ **Certified vs. empirical**: explain when randomized smoothing's sampling cost favors asynchronous auditing or a selective serving path rather than assuming that every certified deployment has the same latency constraint.
- ☐ **Defense budget**: compare input detection with intrinsic model robustness for a low-risk component using measured clean accuracy, robust accuracy, compute, latency, and coverage rather than assuming a fixed clean-accuracy penalty.

14.7 Data Poisoning Defenses

A recommendation model may retain its overall accuracy while promoting an attacker's chosen item to a small group of users. Defending against that failure requires testing the targeted behavior as well as securing the integrity of the training supply chain. Source permissions set the security boundary, and robustness evaluation measures what happens when poisoned records or updates cross it anyway.

Poisoning changes the data from which a model learns rather than perturbing a fixed model at inference time, and its threat runs from early attacks on support-vector machines (Biggio et al. 2012) to web-scale and generative-model data pipelines (Shan et al. 2023). Microsoft's Tay chatbot, which its operator reported was subverted by a coordinated attack within a day of launch (Lee 2016), shows why any system that learns from live interaction makes every user a potential contributor of training evidence. The engineering question therefore becomes which contributions may enter training and how their influence can be audited.

↩ Data poisoning, the corruption of training data or updates so that the learned model misbehaves later, and its sequence of admission, training, and exploitation are established in section 13.5.3.

A dirty-label attack corrupts the label, so a label audit can catch it. Its counterpart, clean-label poisoning, keeps every label correct and corrupts the features instead, which is why several of the defenses below must look past labels to features and representations. Either form can plant a backdoor through the poisoned records, leaving the trained model to behave normally until an input contains the trigger. Generative-model training corpora are exposed to the same threat in a concept-poisoning form, in which a few hundred perturbed images can redirect what a text-to-image model draws for a prompt. Section 13.5.3 examines these mechanisms, that case, and the paths by which poisoned records enter the pipeline.

Poisoned records can resemble legitimate data, so clean aggregate validation accuracy alone cannot establish that a training corpus is trustworthy, and a targeted attack is designed to hide inside otherwise normal performance. Attacks may originate from internal actors with privileged pipeline access or from external adversaries who exploit crowdsourced collection or open data pipelines that lack integrity checks and lineage tracking. The defense therefore runs as four ordered layers along the data path:

1. **Ingress controls** establish provenance and access limits, dictating which sources may contribute or modify data.
2. **Anomaly detection and sanitization** inspect candidate records before training and filter suspicious submissions.
3. **Robust training objectives** reduce the weight that extreme residuals receive for poisoned samples that evade ingestion filtering, although they do not by themselves bound parameter updates.
4. **Post-training checks** test the trained model's behavior on the poisoning objectives that aggregate accuracy hides before the model is released.

Each layer covers a distinct point along the data path and has its own blind spot, so no single detector can substitute for end-to-end controls. Representation learning sits beneath all four as a foundation rather than a fifth gate: it reduces the model's reliance on brittle features that poisoned instances can exploit, but it cannot prove that the training corpus is uncompromised.

For the recommendation system, the interception layer in figure 14.11 examines contributed ratings before they reach training. Its decisions must preserve provenance so later model failures can be traced back to the accepted records.

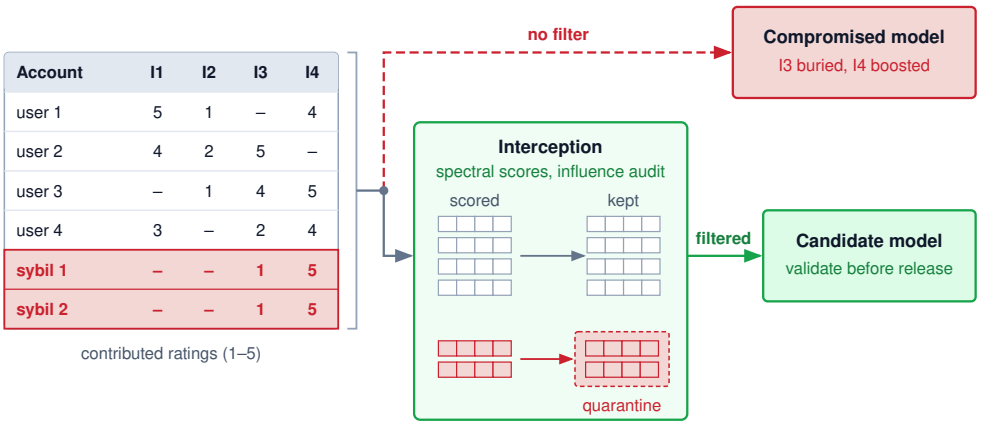

Figure 14.11: Data Poisoning and Interception Workflow: Adversaries inject malicious rows into a user-item rating matrix to manipulate model behavior. A defense layer analyzes contributed data and removes suspected malicious records before training. This can reduce poisoning risk but does not guarantee an uncorrupted model.

14.7.1 Ingress and anomaly controls

If an unverified data source or a compromised contributor account can write directly into a training table, statistical outlier filters will only ever see the attacks that fail to blend in. Poisoning defense therefore begins at the data boundary, before any statistical detector runs. The pipeline must know which sources are allowed to contribute training data, how their authenticity is verified, and which roles can modify accepted records. Strong governance enforces the principle of least privilege,[15] logs data access and modification events, and gives each accepted batch a source identity. These controls do not prove that every record is clean, but they bound the attack surface and give anomaly findings a traceable origin.

> Lighthouse 14.1: Archetype B (DLRM at Scale): Fake profile injection
>
> Archetype B (DLRM at Scale) is especially exposed to **fake profile injection**. Recommendation systems often use online learning to adapt to user trends in real time. An adversary can therefore create a network of "sybil" accounts, fake identities controlled by the same attacker,

[15] **Principle of least privilege**: Articulated by Saltzer and Schroeder (1975), this principle restricts access rights to the minimum necessary for each component. In an ML pipeline, a training job allowed to read unverified sources has an additional path through which adversarial data can enter, so write and read rights become part of the poisoning attack surface. The same principle governs the model registry, the artifact store that maps version identifiers to weight files and controls which alias (for example, `staging`, `production`) resolves to which checkpoint. Promotion to the `production` alias should be restricted to a narrow, audited set of automated evaluation services and named release engineers; if a researcher's credentials are compromised, excluding those credentials from promotion authority blocks that direct path to replacing the live model, while other authorized paths still require scrutiny.

that interact with specific items to artificially boost their popularity or associate them with high-value demographics. This poisoning bypasses traditional firewalls because the malicious interactions look like legitimate user behavior, motivating checks on contributor identity, interaction patterns, and learned representations. Spectral-signature defenses (developed in section 14.7.1.1) search activation patterns for suspicious separation, but their effectiveness depends on whether the poisoned and clean contributions separate in that representation.

After the source boundary, anomaly detection asks whether a candidate training example belongs to the distribution that the pipeline intended to learn from. The cheapest tests compare each record against the bulk distribution: Z-score filtering, Tukey's method, and Mahalanobis distance flag examples whose feature values sit far from normal ranges. These tests are useful because they are fast enough to run at ingestion, but they catch only poisons that look like statistical outliers in raw feature space. For high-dimensional or multimodal data, raw-feature distances are a weak signal: a manipulated image or caption that passes raw-feature checks may separate from clean examples in a pretrained vision or language model's embedding space. That separation depends on the encoder and attack, so it must be measured rather than inferred from the use of a learned representation. Applying Mahalanobis distance or clustering to those representations tests a different hypothesis from raw-feature screening. The pipeline must check whether the resulting outliers correspond to malicious contributions or to legitimate variation it needs to retain.

More coordinated attacks require structure-aware checks. Clustering methods such as K-means, DBSCAN, and hierarchical clustering look for anomalous groups or isolated points rather than single extreme values. The autoencoder ingestion gate in figure 14.12 adds a learned representation to the screening pipeline. The model compresses inputs through a bottleneck layer $\mathbf{z}$, accepting records whose reconstruction error $\mathcal{L}_{\text{rec}}$ falls below threshold τ and quarantining anomalous outliers. Each method has a blind spot. Clean-label poisons need not be statistical outliers, clustering depends on the representation, and autoencoders can reconstruct unfamiliar inputs well or learn an attack pattern from a contaminated reference corpus. Each gate therefore requires evaluation on both poisons and legitimate rare examples.

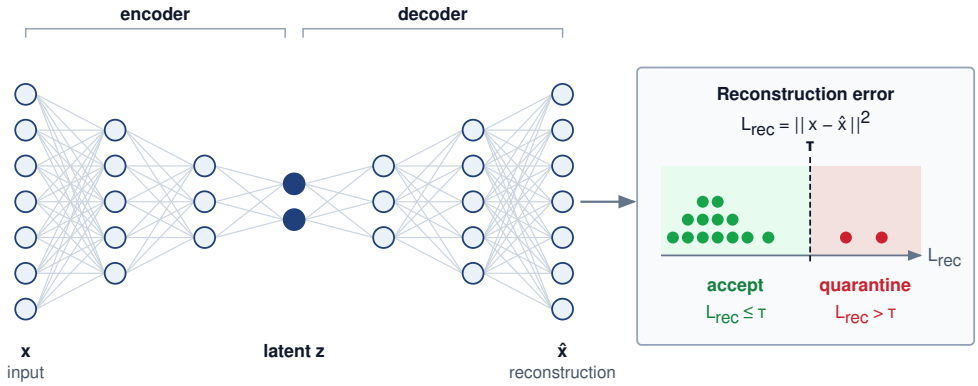

Figure 14.12: Autoencoder Ingestion Gate: An autoencoder trained on trusted data learns to compress inputs through a narrow latent layer $\mathbf{z}$ and reconstruct them. Reconstruction error supplies an anomaly score, not proof that an input is poisoned. The gate accepts records whose reconstruction error $\mathcal{L}_{\text{rec}}$ stays at or below a threshold τ calibrated on held-out clean data and quarantines the rest.

14.7.1.1 Sanitization and preprocessing

Sanitization turns anomaly signals into training-set changes before the poison reaches optimization. A suspicious record can be rejected, quarantined for review, down-weighted during training, or kept with an explicit provenance flag. Routine cleaning still matters: deduplication, missing-

value handling, type checks, range constraints, and cross-field validation remove many low-effort poisoning attempts while improving ordinary data quality.

Data provenance and lineage tracking make those decisions reversible. Each datum needs a record of its source, transformations, validation outcomes, and movement through the pipeline. When a model later shows a poisoning symptom, lineage lets the operator trace model dependencies back to candidate source batches and investigate which contributions caused the behavior. Lineage also lets the operator estimate which trained models consumed confirmed compromised examples and decide whether to remove, relabel, or reweight the affected records.

When suspicious data has already entered the corpus, sanitization moves from source checks to representation checks. Spectral signatures exploit separation between clean and poisoned samples in a learned representation (Tran et al. 2018). In a backdoor attack, the poisoned records carry the target label, so the network must learn two distinct internal pathways to assign it: genuine semantic features for clean examples and the injected trigger feature for poisoned examples originating from other classes. That dual mechanism shifts the mean activation vector of the poisoned subpopulation away from the mean of the clean target-class examples, which creates a variance spike along the direction connecting the two subpopulation means. In the backdoor attacks studied by Tran and colleagues, projecting centered target-class representations onto their top singular vector captured that offset direction, and removing high-scoring samples exposed many poisoned records without requiring prior knowledge of the trigger. The method depends on sufficient mean separation relative to normal within-class variation, so an effective backdoor does not by itself guarantee a usable spectral signature. Filtering therefore requires attack-specific evaluation and checks for discarded clean samples.

Influence functions (Koh and Liang 2017) serve a narrower debugging role. They approximate how infinitesimally increasing a training point z's weight changes the model's loss for a test point z_{test}, without retraining. The fitted parameters are the trained model's weights, and the Hessian is the curvature matrix of the average training loss at those parameters. Its inverse translates a small change in a training point's weight into an approximate change in the fitted parameters. Calculated through inverse-Hessian-vector products, influence $I(z, z_{\text{test}}) \approx -\nabla_\theta \mathcal{L}(z_{\text{test}}, \theta^*)^T \mathbf{H}^{-1} \nabla_\theta \mathcal{L}(z, \theta^*)$, this metric measures local sensitivity to upweighting z. For an average training loss over D examples, removing one point corresponds approximately to a weight change of $-1/D$, reversing the sign and scaling this influence by $1/D$. If a model begins assigning an attacker's target label to one narrow class of inputs, influence functions can rank the training examples that contributed most to that decision. The method avoids forming an explicit inverse Hessian by using iterative approximations, but its results remain approximate in deep nonconvex networks. At modern model scale, influence analysis therefore requires validation against the specific architecture, approximation, and debugging task rather than serving as exact attribution.

14.7.2 Robust training objectives

Under a standard squared-error objective, a single poisoned sample that survives ingestion filtering with a large residual error contributes quadratically to the loss and pulls the gradient far harder than a typical point. To limit the influence of surviving outliers before release evaluation, robust optimization modifies the training objective itself. Robust loss functions such as the Huber loss,[16] the Tukey loss (Beaton and Tukey 1974), and the trimmed mean loss down-weight or ignore the contribution of abnormal instances during training. Regularization techniques (ℓ_1 or ℓ_2 regularization) constrain model complexity and can reduce sensitivity to some poisoned data. At a higher level, minimax[17] and distributionally robust objectives optimize performance over a specified uncertainty set. The objective supplies a formal target; an approximate optimizer does not by itself prove that the trained model attains the target or resists attacks outside that set.

Data augmentation expands training diversity by applying horizontal flips, padded crops, rotations, and color jitter (figure 14.13). The variants preserve the label when the transform respects the class (a horizontal flip need not preserve a digit's label), and they can disrupt some localized poison patterns. Other transformations preserve an attack's trigger, so augmentation does not guarantee dilution. Random subsampling or bootstrap aggregating changes which poisoned records each model encounters; whether combining those models reduces the attack's effect requires evaluation against coordinated poisoning.

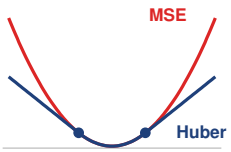

Huber limits a surviving outlier by growing linearly where MSE keeps squaring.

16 **Huber loss**: This piecewise loss transitions from quadratic to linear beyond a chosen residual threshold, bounding its derivative with respect to the residual. The parameter gradient also multiplies that derivative by the model's prediction Jacobian, so Huber loss alone does not bound parameter updates or prevent poisoning. Under MSE, an outlier with 100× normal error contributes 10,000× squared-error loss and 100× residual-gradient magnitude. Under Huber loss, once that residual passes the threshold the loss grows only linearly and the residual derivative stays at the threshold value, so the same outlier pulls no harder than any other point beyond the threshold.

17 **Minimax**: Game-theoretic strategy from Neumann (1928) that minimizes the maximum possible loss. In adversarial robustness, Madry et al. (2018) formulate training as $\min_\theta \max_{\|\delta_{\text{adv}}\| \le \epsilon_{\text{adv}}} \mathcal{L}(f_\theta(x + \delta_{\text{adv}}), y)$: the model learns to minimize loss under a worst-case perturbation within the chosen norm ball. The inner maximization is computationally expensive because it is usually approximated with iterative attacks such as PGD, so PGD-adversarial training is a budgeted robustness choice rather than a free objective change (Bai et al. 2021; Shafahi et al. 2019).

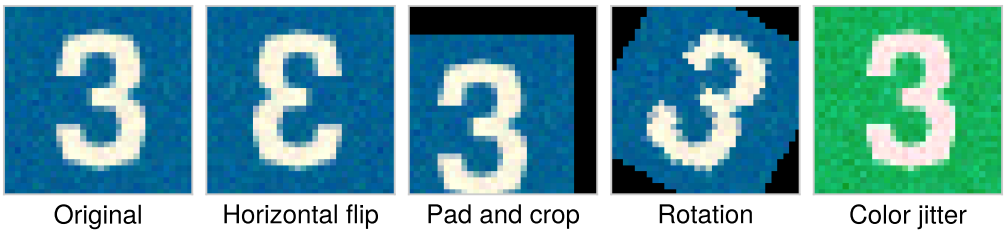

Figure 14.13: Data Augmentation Techniques: Horizontal flips, padded crops, rotations, and color jitter create variants of an existing training example, shown here on an illustrative synthetic digit. When a transformation preserves the task label, it can improve robustness to the represented variation and reduce overfitting; the flipped 3 shows a transformation that does not. The effect on generalization and poisoning resistance requires evaluation.

Distributed and federated training adds a contributor axis to the same threat, and robust aggregation is the training-objective layer's counterpart at the level of contributors. A malicious client need not insert a visible record into a central corpus; it can submit a model update chosen to move the aggregate, implant a backdoor, or exploit sybil identities to gain disproportionate weight. Archetype C (Federated MobileNet) makes this concrete: the server aggregates updates from a large population of smartphone-class clients whose raw data it never sees, so contributor-level checks screen updates before aggregation. Krum (which keeps the one update whose nearest neighbors lie closest to it), coordinate-wise median, and trimmed mean (section 7.3.3) tolerate bounded Byzantine behavior only under their stated assumptions (Blanchard et al. 2017; Yin et al. 2018); high dimensionality and non-i.i.d. client updates can make honest contributors look anomalous. These rules also change the collective, because a median, trimmed mean, or Krum selection needs the individual updates rather than their sum, which raises per-worker traffic over a ring AllReduce.

↩ The traffic cost of gather-based robust aggregation relative to ring AllReduce is established in section 13.6.1.

14.7.3 Post-training behavioral checks

Admitting a poisoned record does not by itself guarantee that training will learn the attacker's behavior, so the last layer tests the trained model rather than its inputs. The evidence that exposes poisoning differs across the four objective classes that section 13.5.3 distinguishes: availability, targeted, backdoor, and subpopulation poisoning. The robustness question is which evidence exposes each class and what that evidence costs (Oprea et al. 2022). Successful availability attacks can lower aggregate validation accuracy, giving broad evaluation a signal to investigate. Targeted attacks, which aim at specific inputs, and subpopulation attacks, which aim at a whole cohort, can leave that aggregate nearly intact, making evaluation on the named targets and cohorts particularly important. Backdoors can also preserve clean-data accuracy while responding to a trigger the evaluator has never seen, as the BadNets pixel-pattern trigger does (Gu et al. 2017). Possible detection strategies include trigger search, canary inputs (probe inputs with known expected outputs, kept out of training), and activation-space inspection when poisoned and clean representations separate (Tran et al. 2018).

For federated and distributed contributors the release gate takes the same form. Held-out behavioral tests on the resulting global model must vary the number and coordination of malicious contributors, their data heterogeneity, participation schedule, and knowledge of the aggregation rule. Global clean accuracy is again insufficient: the gate needs attack success on the trigger or target subpopulation, contributor-level anomaly recall, and the clean-utility cost of robust aggregation. Provenance, participation limits, robust aggregation, and held-out canaries must therefore work together rather than treating one distance rule as a poisoning proof.

The operating rule is layered: authenticate and govern sources before ingestion, sanitize suspicious data before training, limit surviving outliers during optimization, test the trained model on the poisoning objectives before release, and preserve lineage so an incident can be traced back to the compromised source. Data poisoning remains an active research area, but production systems should treat the training corpus as a supply chain rather than a passive dataset.

The four layers act on records, objectives, and the trained model. A representation that does not lean on brittle shortcuts gives a poisoned shortcut less to exploit, and the same property helps

the model survive ordinary deployment variation, so representation learning serves the poisoning defense and the environmental pillar at once. Beyond detecting and correcting shifts after they occur (section 14.4), a complementary approach builds shift-resilient representations from the start, drawing on transfer-learning and domain-adaptation foundations (Pan and Yang 2010).

14.7.4 Self-supervised representations for robustness

Supervised models often overfit to whichever labeled shortcut is easiest to exploit, such as a background texture, an annotation artifact, or a narrow phrasing pattern, and fail when deployment shifts or a poisoning attack targets that cue. Self-supervised learning (SSL) changes the source of supervision so the model learns by solving pretext tasks that require structure in the data itself. A pretraining task that rewards stable structure across views, missing tokens, or masked image patches gives the representation a chance to learn signals that survive more deployment variation.

Contrastive learning methods such as SimCLR (T. Chen et al. 2020) make this idea concrete by pushing different views of the same example toward a shared representation. The model is rewarded for treating a crop, color shift, or augmentation as the same underlying object rather than as a new class-specific cue. **Masked language modeling** in BERT (Devlin et al. 2019) and masked autoencoding in vision (He et al. 2021) use a different route: they hide part of the input and force the model to reconstruct or predict it from context. Both families provide supervision beyond the downstream task labels. Whether the resulting representations transfer better depends on the pretraining data, objective, and deployment shift; removing task labels does not remove shortcut learning.

That transfer benefit is the main systems reason to use SSL in a robustness pipeline. Larger unlabeled corpora expose the model to domains, transformations, and rare cases that would be too expensive to label exhaustively. The same uncurated scale widens the poisoning surface, so SSL pretraining data still needs the ingress and sanitization controls described earlier. In production, SSL usually acts as a foundation rather than a complete defense: pretrain on broad unlabeled data, fine-tune on the supervised task, then apply the adversarial, drift, and poisoning defenses established earlier in this chapter to the task-specific model. Multi-task training can preserve some of this benefit by keeping a self-supervised objective active while the supervised task pulls the representation toward the deployment metric.

The limitation is that SSL is not a robustness guarantee. A contrastive or masked-pretraining objective can still learn brittle shortcuts, and an attacker who understands the pretext task can target those shortcuts directly. The theory explaining when SSL improves robustness remains incomplete, and the compute cost can be substantial because pretraining moves work earlier in the lifecycle. Whether broader representation learning lowers the expected cost of drift, relabeling, and robust fine-tuning enough to justify the added pretraining budget is a per-service calculation. When the model generates open-ended text rather than discrete class labels, the robustness contract itself must expand beyond label stability.

14.8 Robustness of Generative Models

An aligned language model of the Archetype A (GPT-4/Llama-3) class can pass its safety evaluations and still comply with a harmful request once an automatically optimized suffix is appended to the prompt, and suffixes optimized against open models transfer to other models (Zou et al. 2023). That suffix is an adversarial example in token space, found by gradient-guided search over discrete tokens rather than continuous pixels. Generative models also fail without any attacker. Their output is open-ended text rather than a label, so the robustness question changes from whether a label flips under perturbation to whether the generated content preserves factuality, policy constraints, and task intent across prompt variations. A fluent answer can be factually baseless while latency and uptime stay healthy, which makes hallucination the generative form of silent failure.

The evaluation discipline of table 14.1 carries over with three adjustments. First, correctness needs a named reference. A hallucination rate is meaningful only against a stated evidence source, such as the retrieved context, a verified knowledge base, or an expert-labeled answer set, and it counts the generated claims that this reference does not support. The measured rate also depends on the model, the retrieval context, the prompt design, and the sampling temperature, so a reported rate holds only for its prompt set and decoding settings and cannot be carried to another domain

without re-measurement. Second, the perturbation budget becomes a family of prompt variations. Paraphrases, reordered context, and inserted distractors play the role that the ℓ_∞ ball plays for an image classifier, and robustness is reported for each family. Third, the slices are domains and user populations, because an aggregate rate can hide a domain in which the model invents answers, just as the Google Photos failure hid inside aggregate image accuracy.

Detecting a hallucination on an individual request is the per-request competence problem of section 14.4.4 with an open-ended output. Token entropy and log-probability can expose instability in a generation, but neither is a factuality score. A model can assign high probability to a false statement learned from its data, or low probability to a correct but rare statement. Production routing must therefore calibrate these signals against task-specific errors and combine them with evidence checks, retrieval support, or external verification.

The self-consistency technique samples multiple reasoning paths or answers for the same query. Disagreement supplies an instability signal that can route the request for review; agreement supplies weaker evidence because correlated samples can repeat the same misconception or unsupported claim. Sequence-level entropy has the same limitation: it can serve as a calibrated routing feature for a named task, but it does not transform uncertainty into truthfulness.

The signal also has a price. Each additional sample repeats the decode phase for that request, so drawing k samples multiplies its decode work roughly k-fold; only the prompt's prefill can be shared, when the serving system reuses the prompt's cached key-value state. Like the noise samples of randomized smoothing, this cost lands on the compute term of every request it covers, and the same remedy applies. The expensive signal should be spent on the requests whose cheap signals already indicate uncertainty rather than on all traffic. The operational contract must measure error conditional on the routing score and provision the refusal, retrieval, tool, or human path that receives flagged generations.

In retrieval-augmented generation (RAG) architectures, which retrieve documents at inference time and place them in the prompt, robustness depends heavily on the quality of the retrieved context. **Retrieval noise**, the injection of irrelevant or conflicting documents into the prompt, can distract the model, causing it to ignore its internal parametric knowledge and propagate errors from the context. Robust RAG deployments can use re-ranking models, which reorder retrieved documents by relevance to the query, and context verifiers, which check each document before it enters the prompt, to filter out noise before it reaches the generation step. Their evaluation must separate retrieval failure from generation failure. Measuring answer support separately for requests whose retrieved context contains the answer and for requests whose context does not shows whether errors come from the retriever or from a model that ignores or overrides good evidence. The retrieval corpus is also an input an attacker may be able to write to, so a planted document can corrupt retrieval evidence or carry prompt injection, and the ingress controls of section 14.7 apply to it. It becomes a training-poisoning path if the pipeline later admits it into learning.

↩ Reuse of a shared prompt's cached key-value state across requests is established in section 10.7.4.

Language interfaces also carry the input-level attacker into text. Prompt injection, a failure of the boundary between control and data, embeds instructions in user content or retrieved documents so that the model overrides its system prompt. The payload can arrive through a legitimate input channel, so permission to submit content does not make that content trustworthy. External access controls still constrain which data or tools the surrounding system exposes. From the robustness side, evaluation asks whether the content induces behavior outside the intended output contract. It is measured the way adversarial robustness is measured, as an attack success rate over an injection suite matched to the threat model, run by an evaluator who knows the system prompt and filters. Keeping those details secret is not a defense, because an attacker can optimize injection suffixes on other models and those suffixes can transfer to the deployed one.

↩ Retrieval-augmented generation, which keeps domain knowledge in a document index rather than in model weights, is established in section 12.9.6.12.

A common deployment pattern adds **output guardrails**, lightweight classifiers that scan generated text for policy violations or toxicity before the response returns to the user. A guardrail threshold trades the same two quantities as a selective-prediction threshold. Each false positive refuses or escalates a legitimate response, so the false-positive rate multiplied by the legitimate-request rate estimates false refusals or escalations; each false negative is residual risk that reaches the user. Total diverted load also includes correctly flagged harmful responses. The arrival rate at each refusal, retrieval, or human path is the incoming request rate multiplied by the fraction routed to that path under the policy, counting both true and false flags. Choosing the threshold from a curve measured

↩ Prompt injection as a failure of the boundary between control and data is established in section 13.5.5.

on representative traffic, rather than from the classifier's default, allows the team to compare that load with the path's capacity. The guardrail also adds a model inference to every response, so its latency counts against the serving deadline. Because a guardrail is itself a model with its own decision boundary, it carries the same obligation as any other defense and must pass the adaptive evaluation of section 14.6.1.3 against an attacker who knows it is there.

Every routing signal for generative models, from token entropy to a guardrail score, is also computed by the same serving system that the remaining pillar can corrupt.

14.9 System-Level Faults

A preprocessing bug that silently rescales an input feature or an accelerator bit flip that corrupts a stored weight can make a healthy replica emit wrong predictions while triggering the exact same alarms as distribution drift or an adversarial attack. Every defense in the two preceding pillars reads evidence produced by the system it protects: drift monitors read preprocessed features, adversarial detectors read model scores, and poisoning filters read learned representations. The third pillar asks what happens when that machinery is itself wrong. Execution-level detection and recovery belong to fault tolerance. Robustness adds two obligations on top of them: behavioral checks that run on the model's own values and outputs, and a diagnosis that decides whether a behavioral alarm came from a fault at all.

14.9.1 Behavioral checks for silent faults

The exponent bit flip introduced in section 14.3 shows how a single physical upset cascades through neural inference (figure 14.14). In panel (b), flipping bit 30 of an FP32 weight scales its magnitude by 2^{128}, transforming a modest weight w of 0.042 into an extreme outlier w' of about 1.4×10^{37}. Panel (a) traces the affected path through hidden unit h_2 to output logits (y_1, y_2). The corrupted weight is finite and can evade NaN or infinity checks; whether it dominates those outputs depends on the input, activation function, and downstream weights. Targeted bit-flip attacks deliberately select vulnerable weight bits, so their results should not be interpreted as the effect of one random hardware upset. Memory subsystem failures during training can likewise corrupt gradient updates and prevent model convergence. Scale turns one corrupted value into a persistent exposure. Every forward pass of a dense GPT-3-class model reads all 175B of its parameters, so a corrupted weight resident in memory participates in every request that replica serves until something detects it.

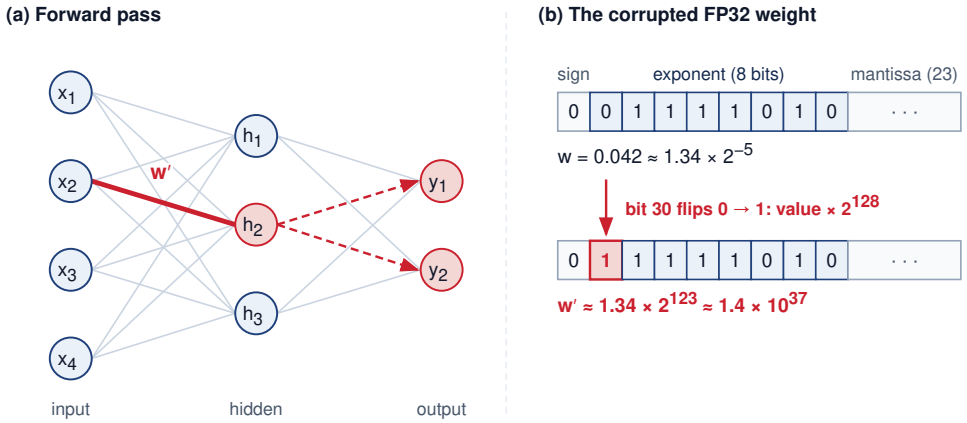

Figure 14.14: Weight Corruption via Bit Flip: (a) One corrupted weight w' drives hidden unit h_2, whose output feeds every downstream logit. (b) For an illustrative trained weight $w = 0.042 \approx 1.34 \times 2^{-5}$, flipping bit 30, the most significant exponent bit of the FP32 word, raises the stored exponent from 2^{-5} to 2^{123}. In general, flipping an exponent bit with place value 2^k scales the value by $2^{\pm 2^k}$; here the factor is 2^{128}, giving $w' \approx 1.4 \times 10^{37}$. The result is still finite, so NaN and infinity checks never fire, yet it can produce extreme downstream activations, depending on the input, activation function, and downstream weights. A range check $|w| \leq W_{\text{max}}$ against the trained weight distribution catches it cheaply.

The cheapest behavioral checks test invariants that a correctly executing model satisfies. A range check $|w| \leq W_{\max}$ against the trained weight distribution catches the exponent flip above at the cost of one comparison per stored value, and the same idea bounds each layer's activations at serving time against the range observed on validation traffic. Where range checks catch extreme values, golden-input canaries replay a fixed set of requests through each replica on a schedule and compare the outputs with stored references, catching a corrupted replica whose liveness and readiness probes still pass. For corruption that stays inside normal ranges and appears only on live inputs, cross-replica comparison runs the same request on two replicas. Because that doubles the work for every request it covers, fleets apply it to sampled traffic, following the check-and-verify pattern used against silent data corruption elsewhere in the fleet. Each check spends along a different C^3 axis: invariant checks add compute on the serving path, canaries and cross-replica comparison add compute and communication, and deciding which replicas and checkpoints remain trustworthy after a check fails is coordination.

↩ Check-and-verify defenses against silent data corruption, and their cost relative to full redundancy, are established in section 7.5.

14.9.2 Distinguishing faults from shift and attack

These checks cover faults that leave a numerical signature. The harder case is a fault whose symptom looks like one of the other pillars. Consider an operator who sees the same surface symptom from all three pillars, falling model accuracy while latency and uptime hold steady. A preprocessing bug that silently rescales a feature can look like covariate shift to a drift monitor, because the feature statistics have genuinely moved. A numerical overflow that corrupts a layer's activations can produce confident misclassifications that look like an adversarial example, because the decision flipped without a visible input cause. The difficulty lies in the disguise, because the cheap pillar-specific detectors fire on the symptom rather than the cause. Three complementary signals guide this investigation: the timing boundary of the change, reproducibility under fixed inputs, and cross-layer correlation with infrastructure telemetry.

The boundary and timing of a change often provide the first diagnostic clue. While environmental shift can evolve gradually or appear abruptly with real-world events, a shift synchronized with a deployment, dependency bump, or schema migration raises the pipeline-fault hypothesis. This suspicion strengthens when multiple features shift simultaneously in directions that the underlying physical process cannot explain. Timing alone narrows the search space without proving causation.

Reproducibility under fixed inputs helps locate the failure boundary. Replaying golden inputs through version-pinned preprocessing pipelines and model artifacts can expose the stage where outputs begin to differ. A deterministic software defect fails at the same stage on every replay. A transient upset in the computation usually vanishes on replay, whereas a corrupted value resident in memory, such as a flipped weight, reproduces on the same replica and disappears only on a second device or after the weights are reloaded. Race conditions and marginal hardware fail intermittently, so isolating them takes repeated runs or a second device. Stochastic execution also requires controlling random state or comparing repeated runs. Because shifted or adversarial inputs can reproduce the same error under fixed execution conditions, replay alone cannot settle the diagnosis; operators must compare raw and preprocessed tensors, artifact versions, and population-level statistics alongside the replay output.

Finally, cross-layer correlation links model behavior to surrounding infrastructure telemetry. Query-rate spikes or unusual input patterns can support an attack hypothesis, while deployment releases, ECC memory counters, SDC checkers, or memory-pressure alarms can support a system-fault hypothesis. Neither pattern is necessary or sufficient. An attack can remain within normal traffic patterns, and silent corruption can escape the available counters. Correlation guides the next diagnostic test rather than establishing the cause.

Example 14.2: A drift alert that was a pipeline fault

Scenario: A credit-risk service raises drift alerts on three income-related features on the morning after a routine dependency update, and the approval rate for applicants in the affected income bands shifts sharply. Latency, error rates, and uptime are unchanged, and every request that the on-call engineer retries returns a score without error.

Diagnosis: The three signals point the same way. The alerts begin at the release boundary, and the three features moved together in a direction that no change in applicant behavior explains. Replaying a stored set of golden requests through the previous and current pipeline versions produces identical raw inputs but preprocessed tensors that differ by a constant factor on exactly those features, because the updated library changed a default normalization. Deployment telemetry correlates with the change; query-rate and hardware counters do not. The failure is a software fault disguised as covariate shift. Retraining on the corrupted features would have taught the model the bug, and the successful retries proved only that the service answered, not that its answers passed acceptance.

Systems lesson: A replay harness that pins artifact versions and compares tensors at each pipeline stage turns an ambiguous drift alert into an attributable fault. Recovery is complete only when the replayed outputs satisfy the model's acceptance tests again, which is a behavioral check that a successful retry or a green health dashboard cannot supply.

The paired replay in figure 14.15 isolates the first changed stage in this example by holding the raw requests and model artifact fixed.

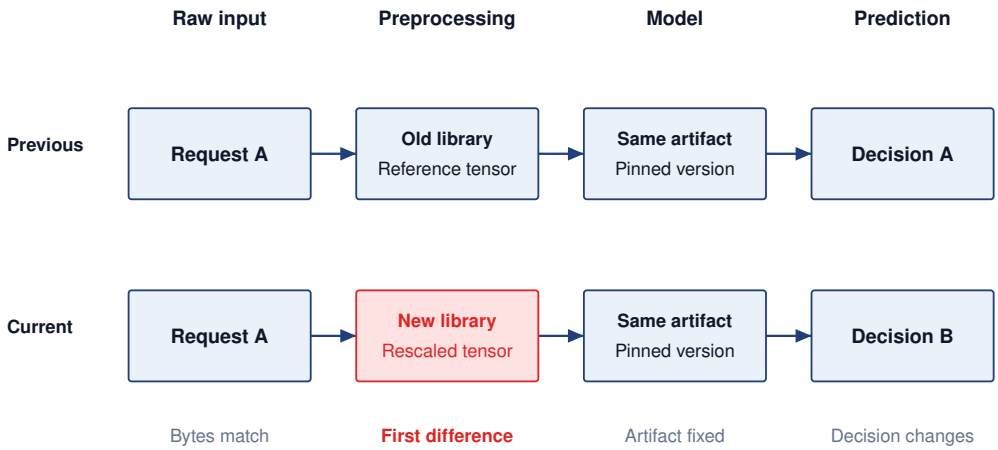

Figure 14.15: Locating a Pipeline Fault by Replay: Previous and current pipelines receive identical stored request bytes and use the same model artifact. Their preprocessing libraries differ, and the first differing tensor appears immediately after preprocessing. The changed decision is downstream evidence, while the earlier tensor difference identifies the boundary to investigate. This controlled example supports a pipeline-fault diagnosis; replay alone does not distinguish every fault, attack, or population shift.

Comparing only the final scores would expose a behavioral change without locating it. Comparing intermediate tensors directs the acceptance test and rollback toward preprocessing before a retraining decision.

The diagnostic discipline is therefore to read the system-level evidence before accepting the drift or attack hypothesis the surface symptom suggests, because each pillar suggests a different intervention and more than one cause can be active at once. Skipping that diagnosis, or mistaking a passed clean-data or single-norm test for comprehensive protection, produces the recurring engineering pitfalls examined next.

14.10 Fallacies and Pitfalls

Robustness spans environmental shifts, input-level attacks, and system-level faults. Each threat domain introduces misconceptions that lead to inadequate defenses or misallocated engineering resources.

Fallacy: *Adversarial examples are an academic curiosity with no real-world impact.*

Published physical-world attacks, such as adversarial patches on clothing or stickers on stop signs, have fooled evaluated vision models under controlled patch and road-sign settings without digital access. Defense requires evaluation under the relevant physical conditions. Physical-world data augmentation, patch-aware training, and serving-time detection are candidate interventions, each with its own coverage and cost. PGD-style adversarial training on digital images adds work with each attack step, but its results do not establish resistance to physical attacks. Ordinary digital test sets can miss these deployment failure modes.

Pitfall: *Treating test-set success as robustness proof.*

A test set drawn from the same i.i.d. distribution as training data measures performance in that population; it does not establish resilience to a different deployment distribution or to a critical slice inside the same one. The Google Photos labeling failure shows how a harmful error on one population slice can hide inside aggregate accuracy, and the night-driving and clinic-imaging examples of distribution shift fail on populations the test set underrepresents. Robustness evidence needs held-out shifts, named slices, and the other cells of table 14.1.

Fallacy: *Distribution shift monitoring is optional after deployment.*

Models often degrade silently, maintaining high confidence scores even as predictive performance falls due to drift. Monitoring metrics like PSI can surface population shifts before accuracy falls below service level agreement (SLA) thresholds, enabling proactive intervention when the monitored features are predictive of downstream performance.

Pitfall: *Ignoring poisoning defenses unless the attacker controls the training pipeline.*

Clean-label poisoning compromises models by injecting correctly labeled but manipulated samples into public datasets or scraped data sources, requiring no access to internal code. Low-dose attacks can still target a particular prediction or implant a backdoor; the size of the poisoned slice and the apparent correctness of its labels do not establish integrity.

Fallacy: *Average accuracy is sufficient for robustness measurement.*

High average accuracy often masks fragility. A model that is accurate on nearly all clean test inputs can still fail on every input in a critical slice once an attacker applies a small, targeted perturbation (Madry et al. 2018). Reliable evaluation requires evidence for the claimed failure class: slice-level performance under shifts, empirical attack accuracy under a stated budget, or a certified radius where a formal guarantee is available. A finite attack search supplies empirical evidence rather than exact worst-case accuracy.

Pitfall: *Treating adversarial training as a complete robustness solution.*

Adversarial training does not establish universal robustness. A model trained against ℓ_∞ attacks needs separate evaluation for ℓ_2 or geometric attacks. Perturbation sets can overlap, so some resistance can carry over to other norms, but the training threat model alone does not establish the extent of that transfer.

Fallacy: *Software faults do not matter when the model is correct.*

Focusing solely on algorithmic robustness ignores the reality that software bugs in data pipelines and serving infrastructure are a major cause of ML failures. Pipeline and data issues, not model architecture or adversarial attacks alone, can sit at the center of failures. Checksums can detect changed bytes, while behavioral validation must also catch incorrect transformations of intact data. A model hardened against adversarial perturbations can still fail when a preprocessing bug supplies validly stored but wrongly computed inputs.

Pitfall: *Hardening the model while leaving pipeline checks untested.*

Robustness work can concentrate on adversarial training, certified radii, or poisoning defenses while leaving schema validation, preprocessing parity, feature freshness, and rollback drills under-tested. That imbalance creates a system that resists one class of attack but fails on ordinary operational faults. In the credit-risk example of section 14.9.2, a changed library default would have been caught at release time by a golden-request replay comparing preprocessed tensors across pipeline versions.

Without that check, the fault surfaced as a drift alert, and retraining on the corrupted features would have taught the model the bug. A robust deployment validates the model and the pipeline together and exercises the replay harness and rollback path before an incident needs them.

14.11 Summary

Robust AI makes prediction quality an explicit engineering obligation alongside uptime and latency. Environmental shifts, adversarial manipulation, poisoning, and system faults can leave a service apparently healthy while its outputs become unreliable. The chapter's diagnostic framework reads release timing, fixed-input replay, and infrastructure telemetry before accepting a drift or attack hypothesis, then connects each failure class to its own evidence and response. Fleet recovery mechanisms must therefore be paired with behavioral acceptance tests.

Inside the training step, robustness is paid for in the terms of the Fleet Law (equation 1.3). Multi-step PGD adds gradient work in T_{compute}; at a fixed worker count N, the wall-clock increase depends on implementation, reuse, and the other terms in the step. Whether it also reduces clean accuracy, and by how much, depends on the model, data, perturbation set, and training recipe. Robust aggregation for distributed or federated contributors replaces a reducible mean with a median, trimmed mean, or Krum selection, which moves cost into $T_{\text{comm}}(N)$ because the individual updates must be gathered rather than summed.

Outside the training step the same C^3 axes set the price. Randomized smoothing and self-consistency sampling spend compute on every request they cover. Invariant checks and golden-input canaries spend compute and communication to catch faults that health probes miss. Poisoning sanitization spends compute and storage bandwidth on extra passes over the training corpus. Drift monitoring spends communication, shipping reference statistics and serving telemetry between machines. Release gates, abstention thresholds, and escalation ownership spend coordination, because validated model versions must be adopted consistently across the fleet.

Conformal routing touches all three axes: computing prediction sets adds bounded compute, sending requests with large sets to a stronger model adds communication, and routing abstentions to reviewers adds coordination. If prediction sets grow and trigger more deferrals, human-review or fallback capacity can become the bottleneck. Distribution shift can also invalidate coverage without enlarging the sets, so set size is a load signal rather than a complete drift detector. A coverage claim is operationally useful only when the deployment also provisions the path that consumes uncertain cases. Robustness engineering measures these workload-specific costs against the reduction in silent risk where failure consequences are highest.

Key Takeaways: Silent failure is the real threat

- **Silence is the failure mode**: Robustness exists because drift, adversarial perturbations, poisoning, and numerical faults often preserve uptime and latency while corrupting predictions. Production systems need monitors that surface competence loss before user complaints or downstream business metrics reveal it.
- **Diagnose before defending**: A software fault can look like covariate shift, a poisoned sample like a rare outlier, and an adversarial input like natural variation. Read the release timing, fixed-input replay, and infrastructure telemetry before spending a defense budget on the hypothesis the surface symptom suggests, then combine ingress validation, training defenses, uncertainty signals, and output verification.
- **Robustness is bought explicitly**: Strong adversarial training (principle 24) can reduce clean accuracy and multiply training compute, while certified defenses and uncertainty sampling add serving cost. Price each effect against a like-for-like baseline before deciding how much resilience the failure consequence justifies, and claim robustness only for the evaluated attacker access and perturbation budget.
- **Poisoning defense follows the data path**: Poisoned records can leave clean aggregate accuracy intact, so no single detector suffices. Layer source controls, sanitization, robust objectives, and post-training behavioral checks, and keep lineage so every admission decision can be traced and reversed.

- **Drift alerts begin an investigation**: Calibrated distance measures identify changes that need review but can stay silent under concept drift that leaves the input distribution unchanged, so outcome evidence must separate a change in label frequencies from a change in the input-to-label relationship. The shift type and root-cause analysis then determine whether the response should be continued monitoring, data repair, recalibration, retraining, or fallback.
- **Per-request contracts have operating conditions**: Raising a selective-prediction threshold lowers risk only if the score ranks requests by error risk, and conformal prediction provides finite-sample marginal coverage under exchangeability, not per-input or per-subgroup certainty. Version and refresh the calibration evidence, and provision fallback capacity for rejected requests and large prediction sets.
- **Generative reliability is semantic**: Large language model hallucinations are confidently fluent failures rather than simple label mistakes. Entropy and self-consistency are routing signals, not factuality guarantees; grounding checks, external verification, and escalation policies must bound what the system is allowed to assert.

No finite evaluation establishes robustness to every shock. A deployable claim must bind behavioral evidence to the evaluated artifact and operating conditions, then provision the detection and fallback paths without violating service-level objectives (SLOs). The release record in section 12.5.8 preserves that link so a changed model or workload cannot silently inherit obsolete evidence. Governance cites the resulting record as evidence in an assurance case, a versioned argument that a model is fit for a stated use. Governance uses that case to assign acceptable residual risk and decision authority; a robustness score alone cannot make those decisions.

↪ The assurance lifecycle that turns such evidence into launch and stop decisions is developed in section 16.2.

What's Next: From resilience to sustainability

Even when a fleet is secure and robust, it can still be bounded by the electrical power its accelerators demand. Every defense in this chapter is paid for in compute. Adversarial training multiplies gradient work, randomized smoothing and self-consistency sampling repeat inference, and replay harnesses and redundant checks run beside the serving path. That compute draws electricity and rejects heat, and a facility's power and cooling capacity bounds how much of it the fleet can run. The next chapter treats those limits as design constraints, beginning with the energy ceiling on fleet capacity and extending to operational and embodied carbon and carbon-aware operation.

↪ Facility power and cooling limits, lifecycle carbon accounting, and carbon-aware scheduling are developed in chapter 15.

15

Sustainable AI

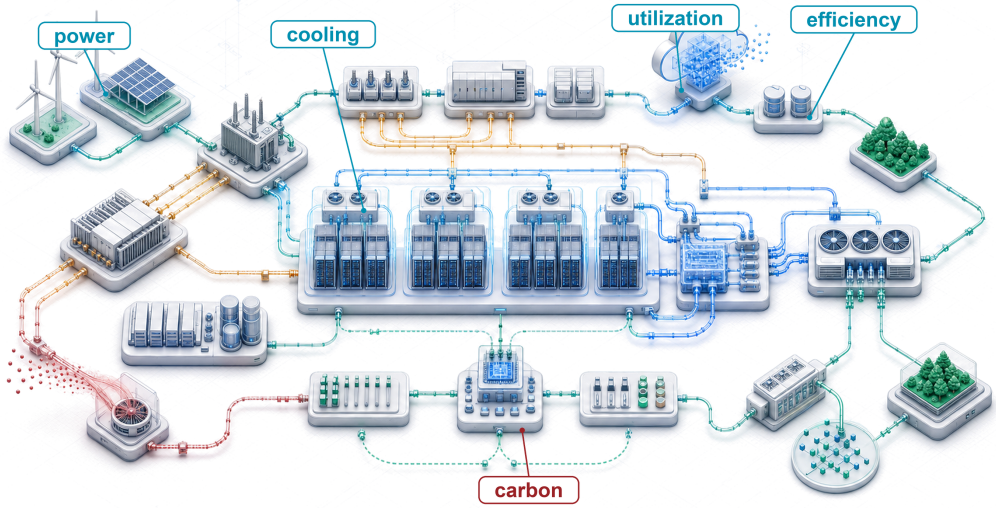

15.1 The Energy Ceiling
15.2 The Scale of Environmental Impact
15.3 Exponential Growth vs. Physical Constraints
15.4 Energy Measurement and Modeling
15.5 Carbon Footprint Calculation
15.6 Data Center Energy and Resource Consumption
15.7 Training vs. Inference Energy Analysis
15.8 Hardware Lifecycle and E-Waste
15.9 Mitigation Strategies
15.10 Policy, Regulation, and the Path Forward
15.11 Fallacies and Pitfalls
15.12 Summary

Purpose

Why do power limits determine what machine learning systems can exist, not just what they cost to operate?

Power is an operational expense and a physical constraint on what can be built. A data center has a power budget determined by electrical infrastructure and cooling capacity; a training run that exceeds it cannot proceed at the intended scale regardless of its financial budget. Adding accelerators only helps if the site can supply their electricity and remove their heat. A design that runs efficiently on one machine can therefore become infeasible when replicated across a fleet. The consequences reach beyond a delayed training run. Power reserved for training may be unavailable to serving, while equipment kept ready for failures consumes resources even when it performs no useful work. Moving work to another location can relieve one constraint while changing its emissions, access to data, or ability to meet deadlines. Efficiency improves these choices, but lower energy per request does not ensure lower total consumption when demand grows. Nor does a clean electricity supply erase the environmental cost of manufacturing and replacing equipment. Engineers must judge feasible designs over their useful lives and account for the communities sharing their water and electricity supplies. Sustainability therefore belongs alongside accuracy, latency, and reliability in system design: a fleet must deliver useful service within resource limits that its performance targets cannot override. In C^3 terms, compute, communication, and coordination draw on a shared physical budget, so their combined resource demand determines whether the fleet can operate at all.

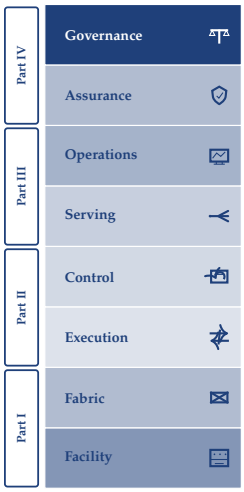

Learning Objectives

- Explain why facility power and cooling limits constrain which ML systems can physically exist
- Analyze how scaling-driven compute growth, hardware efficiency, and demand rebound shape fleet resource demand
- Calculate operational, embodied, and lifecycle carbon from energy use, power usage effectiveness, manufacturing, and grid intensity
- Analyze how computation, memory movement, power delivery, cooling, and water availability constrain sustainable fleet operation
- Compare training, inference, edge, and federated workloads by energy, carbon, and battery budgets
- Design mitigation strategies across algorithms, infrastructure, scheduling, hardware lifetime, and carbon-aware operations
- Evaluate sustainability claims against a measurement contract that fixes functional unit, system boundary, and carbon signal

15.1 The Energy Ceiling

A model can be secure against adversaries and robust under distribution shift yet still fail if the available grid cannot power its accelerators or the facility cannot cool them. **Sustainability** begins at that boundary, where accuracy, latency, and reliability are no longer enough because the system must also fit within power, cooling, water, carbon, and hardware lifetime budgets. Every operation a model performs is paid for in energy, and power density sets the ceiling on data-center computational capacity.[1]

When an engineer optimizes a database query to save 100 ms, that is ordinary performance tuning. At fleet scale, the same saving repeated billions of times per day can reduce electrical demand, cooling load, and carbon emissions if it reduces measured energy per completed query (Lacoste et al. 2019).

The power ceiling makes sustainability a design constraint, not a reporting category. A system that exceeds its operating envelope has failed operationally in the same sense as one that crashes, because it cannot be deployed at the intended scale. The same engineering discipline that budgets memory, bandwidth, and fault tolerance must also budget joules, cooling load, embodied carbon, and hardware replacement cycles.

Physical feasibility establishes which fleets can run. Sustainability then compares the environmental cost of feasible designs over their useful lives. The account starts from workload energy at a declared measurement boundary and extends to carbon, water, and manufacturing. The dominant lifecycle term then decides among efficiency, placement, and hardware reuse. Communication and coordination belong in that account alongside compute, because transfers, synchronization stalls, and recovery keep equipment powered even when useful computation stops. A distributed step costs its average fleet power multiplied by the step time that the Fleet Law predicts, so every second of exposed communication or synchronization is also a block of joules. Equation 15.13 writes that account out where the chapter turns to distributed training.

Contemporary machine learning applications operate at industrial scale, with individual campuses drawing power at the level of heavy industrial plants. Training a single large AI model can consume as much electricity as roughly 120 US homes do in an entire year. The scaling laws that motivate fleet-scale training (principle 9) price each further loss reduction in additional compute, data, and parameters, so compute demand has grown faster than hardware efficiency (Sevilla et al. 2022). That scaling-driven growth is the first way demand outruns efficiency, and sustainable AI treats it as a first-class systems problem rather than an externality to be managed after deployment.

[1] **Joule energy scale**: An A100 modeled at its rated thermal design power consumes 400 J every second (400 W); actual workload power requires measurement. Large-model training has no single defensible joule range: model size, hardware count, realized power, utilization, and run duration can each shift total energy by orders of magnitude. The engineering lesson lies in the conversion from repeated per-operation costs to facility energy, so this chapter reports measured or modeled workload energy rather than assigning every large run to one fixed range.

↩ The power-law scaling of loss with compute, data, and parameters, and the compute-optimal allocation that follows from it, are established in section 1.3.

Definition 15.1: Sustainable AI

Sustainable AI is the systems engineering practice of measuring and optimizing the full environmental cost of ML systems (energy, water, and embodied carbon across training, inference, and hardware manufacturing) and incorporating those costs as explicit constraints in architecture decisions alongside performance and accuracy objectives (Wynsberghe 2021; Lannelongue et al. 2021; Henderson et al. 2020).

1. **Significance**: Training GPT-3 consumed approximately 1,287 MWh of energy (Patterson et al. 2021), equivalent to roughly 120 US household-years of electricity. The same lifecycle logic applies after training: pretrained models can often be adapted with far less work than training from scratch, and high-volume inference can eventually dominate total energy. Sustainable AI therefore tracks both one-time training runs and recurring serving workloads rather than treating model training as the whole footprint.
2. **Distinction**: Unlike corporate sustainability reporting (which aggregates energy usage into annual CO_2 disclosures), sustainable AI engineering operates at the individual workload level. It selects hardware by FLOP/s per watt, schedules training during periods of high renewable availability, and chooses model architectures that minimize inference FLOPs rather than maximizing accuracy alone.
3. **Common pitfall**: A frequent misconception is that switching to renewable energy solves the sustainability problem. For hardware-intensive ML, embodied carbon (the carbon emitted manufacturing the chips, servers, and cooling equipment before they ever run a training job) can rival operational carbon on clean grids or at low utilization. For battery-powered consumer devices, manufacturing can account for most of the lifecycle footprint (Gupta et al. 2022). Hardware longevity and utilization rate are therefore as important as energy source.

The environmental impact of AI systems spans the complete lifecycle: from semiconductor manufacturing and data center construction to model training, inference deployment, and electronic waste (Wynsberghe 2021; Gupta et al. 2022; Luccioni et al. 2023). Treating this full lifecycle as an engineering problem rather than a corporate responsibility exercise transforms sustainability from a vague objective into a measurable engineering requirement. Optimizing this footprint requires first grounding intuition in the raw physical energy of a single large-model training run.

Any such calculation first has to fix what the energy measurement includes. This is the accounting discipline that the energy accounting boundary (equation 1.5) imposes on fleet comparisons, applied here to a single run. Accelerator energy excludes host processors, host memory, networking, and storage. IT energy includes those supporting systems; facility energy additionally includes cooling and facility-side power losses.

Power usage effectiveness (PUE) is the ratio of facility energy to IT energy over the same interval (equation 15.2); section 15.4.1.3 gives its formal definition and what it leaves out. A run's fleet energy productivity is the useful FLOPs it completes per joule at a stated boundary. Dividing useful training FLOPs by that productivity gives energy at the boundary over which the productivity was measured. If that boundary is IT equipment, multiplying by PUE gives facility energy; if it already includes the facility, applying PUE again double counts overhead. The opening calculations use one completed training run as their unit of work and account for operational electricity. Manufacturing and other lifecycle terms enter separately.

Checkpoint 15.1: The energy of intelligence

A 175B parameter model requires approximately 3.14×10^{23} FLOPs to train. Assuming a data center PUE of 1.1 and a realized training efficiency of 50 GFLOP/J measured across the IT equipment, including support systems but excluding facility overhead:

- ☐ **Training energy**: calculate the total energy consumption in megawatt-hours (MWh).

- ☐ **Household-years**: convert the training run's energy into household-years using the average US household consumption of 10.7 MWh per year.
- ☐ **Reported value**: compare the calculated estimate with the 1,287 MWh reported for GPT-3 in the definition above, and identify which assumption (the realized efficiency or the PUE) would have to change to reconcile the two.
- ☐ **Metric boundary**: assess whether household-years captures the true environmental cost by separating energy consumption (MWh) from carbon intensity, the grams of CO_2 the grid emits per kilowatt-hour delivered (g CO_2/kWh).

That household-year figure is the stake; it is also only the electricity entering one building for one run. Pricing its environmental impact requires converting those megawatt-hours into carbon emissions across regional power grids, placing single-run electricity draw inside global data-center demand, and accounting for the water, manufacturing, and hardware disposal burdens that accompany the electricity bill.

15.2 The Scale of Environmental Impact

A training run measured in megawatt-hours is hard to reason about, because few engineers carry an intuition for what a megawatt-hour of grid electricity costs the atmosphere. The introduction quoted GPT-3's training emissions as one of the costs that scaling imposes (section 1.3.4). The derivation behind that figure converts the run's energy into emissions with a grid factor and then into a familiar anchor, the CO_2 of a trans-Atlantic flight, which turns an abstract number into one with physical stakes.

Napkin Math 15.1: The carbon cost of training

Problem: A team trains a large model (GPT-3 size), treating 1,287 MWh as facility electricity for this comparison. How much CO_2 is emitted under the stated grid factor, and how does that compare to a trans-Atlantic flight?

Math:

1. **Energy**: Training energy consumption is 1,287 MWh = 1,287,000 kWh.
2. Carbon intensity (US average): $\approx$ 0.429 kg/kWh.
3. Total Emissions: Training emissions are 1,287,000 kWh $\times$ 0.429 kg/kWh $\approx$ 552,123 kg.
4. Comparison:
   - One passenger, NY to London (round trip): $\approx$ 1000 kg.
   - Ratio: 552,123 kg / 1000 kg = 552.1.

Systems insight: Under these assumptions, a single training run emits as much carbon as hundreds of trans-Atlantic passenger round trips. The calculation holds facility energy fixed, so changing the grid factor changes the attributed emissions. The same energy on a hydro-heavy grid would emit a small fraction as much, the placement lever that notebook 15.2 quantifies.

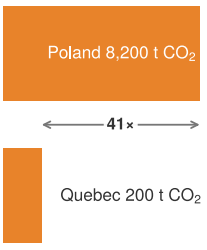

The same job's carbon emissions depend on where it runs.

The same run's emissions depend on where its electricity comes from. For a fixed workload, the grid emissions factor and the facility overhead together determine the attributed carbon, and the placement calculations separate those effects, starting with the grid factor alone.

The introduction previewed this geographic factor as a spread of roughly 41× between hydro- and coal-dominated grids. Notebook 15.2 derives that spread by holding facility energy constant and varying only the grid factor.

Napkin Math 15.2: The geography of carbon

Problem: A team is choosing a data center for a 10,000 MWh training run, comparing Site A in Quebec (hydropower, 20 g/kWh CO_2) against Site B in Poland (coal-heavy, 820 g/kWh CO_2). Both factors are rounded annual grid averages on the same emissions boundary, the values figure 15.7 later applies across nine regions; they are not hourly or marginal signals. How does the location affect a model's carbon footprint?

Math: Carbon = Energy × Grid Intensity.

1. Site A emissions: 10,000,000 kWh × 20 g/kWh = 200,000,000 g = 200 t CO_2.
2. Site B emissions: 10,000,000 kWh × 820 g/kWh = 8,200,000,000 g = 8200 t CO_2.
3. Ratio: 8200 t / 200 t = 41× difference.

Systems insight: The Quebec-versus-Poland pair gives a 41× difference in attributed emissions at the same facility energy. This comparison isolates the grid factor; it does not rank placement against an unmeasured algorithmic change. A fleet decision must compare feasible sites and schedules using both workload energy and the emissions associated with supplying it.

The ratio holds only at equal facility electricity and these annual-average factors. A second gain from waiting for cleaner hours, the temporal lever of section 15.9.4, therefore cannot simply be multiplied onto it; the joint schedule must be evaluated as one decision.

The GPT-3 household-year calculation above anchors the workload scale. The fleet-scale consequence is larger. The International Energy Agency (IEA) projects global data-center electricity consumption to reach about 945 TWh by 2030, just under 3 percent of global electricity demand, with AI-accelerated servers driving much of the growth (International Energy Agency 2025).[2] Historical studies document rapid growth in training compute (Sevilla et al. 2022), but a compute multiplier is not an energy multiplier without a matching efficiency measurement.

Beyond direct energy consumption, AI systems drive environmental impact through hardware manufacturing and resource consumption. Producing processors and their supporting equipment requires materials and industrial processes whose impacts precede the first training run.[3] At the other end of the hardware's life, retired equipment joins the electronic-waste stream whose scale section 15.8.5 quantifies. Rising workload requirements can motivate accelerator replacement before physical failure; whether replacement creates waste depends on the equipment's reuse and disposal path.[4]

These environmental challenges are part of the energy ceiling, not an add-on to it, because the communities that host land, water, grid capacity, and waste streams also carry part of the system cost. These scale calculations turn sustainability into an allocation problem. AI progress creates benefits in one part of the system while assigning electricity, water, land, and disposal costs to another, so environmental responsibility must be treated as part of systems design rather than as postdeployment reporting (section 15.10.4 examines how site selection distributes these burdens across communities). A static snapshot of today's footprint still understates the engineering challenge, however, because the computational demand driving these electricity, water, and hardware budgets is compounding far faster than hardware efficiency can absorb it.

15.3 Exponential Growth vs. Physical Constraints

Between 2012 and 2023, the training compute of frontier models expanded from roughly 4.7×10^{17} FLOPs for AlexNet to 2×10^{25} for GPT-4-class systems, an increase of about 45 million times in 11 years that outpaced hardware efficiency gains by orders of magnitude. Sevilla et al. (2022) estimate that the training compute of notable machine learning models doubled roughly every six months after 2010, far faster than the two-year cadence historically associated with transistor scaling. Extending the introductory series (figure 1.1) across 468 notable runs through 2026 (Epoch AI 2026c), figure 15.1 shows that the 69 runs flagged as frontier at release grew about 4.6× per year, doubling every 5.5 months and carrying the post-2012 span past nine orders of magnitude to 1×10^{27} FLOPs. Larger models, longer training runs, and rising serving volume each add to that total work through

2 | **Data center industrial scale**: The IEA projection is an electricity-demand metric, not directly comparable to aviation or cement shares of global emissions. AI infrastructure competes with other loads for deliverable power; expansion requires available capacity, additional supply, or a reduction or shift in competing demand, not necessarily new generation within the same region.

3 | **GPU manufacturing embodied carbon**: NVIDIA's HGX H100 product carbon footprint reports 1,312 kg CO_2e (CO_2-equivalent, which expresses other greenhouse gases as the mass of CO_2 with the same warming effect) for an eight-H100 baseboard, or roughly 164 kg CO_2e per H100 if allocated evenly across the eight GPUs (NVIDIA Corporation 2025). Advanced-node manufacturing also consumes ultrapure water, specialty gases, and chemicals that an energy-only account never sees, so a per-GPU carbon figure understates the full manufacturing footprint.

4 | **AI hardware e-waste**: Retiring an accelerator does not necessarily create waste: equipment may be redeployed, resold, or recycled. Estimate disposal impacts from the actual equipment inventory, material composition, and end-of-life route. Extending useful service life can defer manufacturing and disposal, while certified handling limits exposure to hazardous materials in discarded electronics.

a different mechanism. Energy per operation can fall while total fleet energy rises; the outcome depends on how much useful work the fleet performs and how efficiently it performs that work.

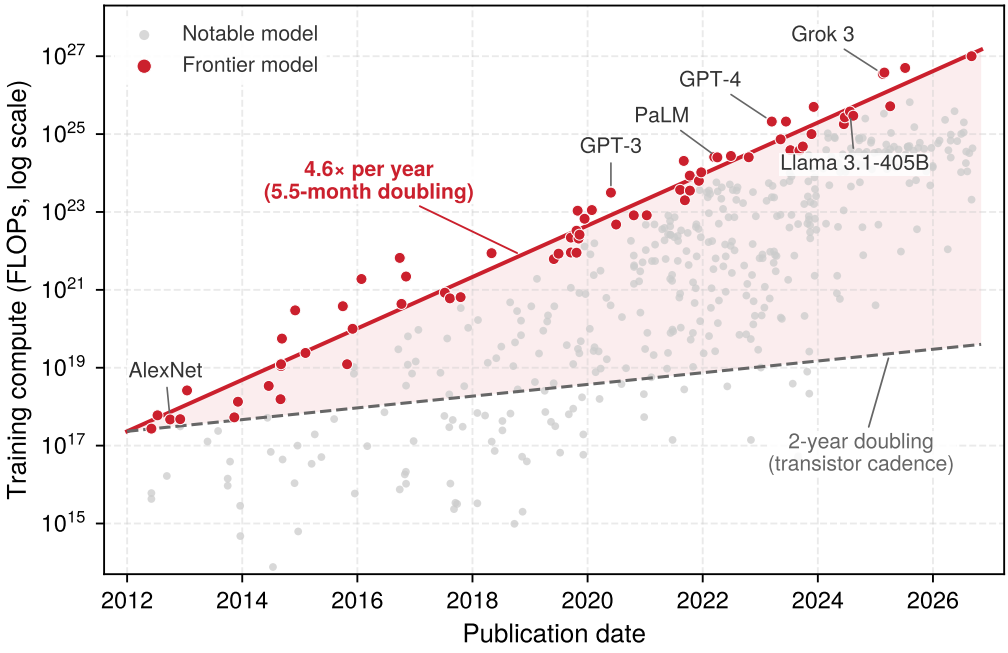

Figure 15.1: Training Compute of Notable Models: Training compute of 468 notable models published from 2012 through 2026, on a logarithmic axis. Red circles mark the 69 runs flagged as frontier at release; their least-squares trend grows about 4.6× per year (95 percent confidence interval 4.3× to 4.9×), doubling every 5.5 months. From the same 2012 origin, the dashed line doubles every two years, close to Moore's law, and the shaded gap shows how rapidly frontier demand pulls away from transistor-cadence scaling. Because the vertical axis counts operations rather than joules, converting this trajectory into energy demand requires the realized efficiency of each run at a declared measurement boundary. Data from Epoch AI (2026c), licensed under CC BY 4.0; many values are Epoch AI estimates, and the post-2010 era framing follows Sevilla et al. (2022).

Historically, semiconductor scaling combined greater transistor density with lower switching energy. Moore's law described the growth in transistor count, while Dennard scaling explained how reducing dimensions and voltage together could preserve power density. Transistor count alone did not guarantee better energy efficiency. Once voltage scaling encountered leakage and other device limits, a denser chip could no longer assume proportional power savings.[5] Further efficiency gains require improvements in devices, architecture, and the work performed by the application.

At a fixed accounting boundary, energy depends on total work divided by realized work per joule. Improvements in execution efficiency reduce IT energy, and PUE then adds the facility overhead on top of that IT energy. Algorithmic changes can also reduce the work required to meet a quality target. These effects must be measured together rather than inferred from a historical compute curve.

Per-job efficiency and total demand can move in opposite directions. Fleet energy is the energy per completed job multiplied by the number of completed jobs, so a fleet forecast must state both factors, including the workload mix. Cheaper jobs can raise the job count faster than they lower the energy of each one. If energy per task falls tenfold while task volume rises a hundredfold, fleet energy rises tenfold. This rebound is a second way demand can outrun efficiency, distinct from the scaling-driven growth described at the start of the chapter. Scaling-driven growth concerns how much compute each frontier model requires as scaling laws push quality upward; rebound concerns how many tasks run once each task becomes cheaper. The mitigation strategies of section 15.9 treat rebound under its historical name, the Jevons paradox.

[5] **Dennard scaling**: Robert Dennard observed in 1974 that smaller transistors could operate at constant power density by reducing voltage proportionally. This scaling pattern ended around 2005 when leakage current made further voltage reduction impractical. The discussion of Dennard scaling in section 2.6.1 drew the facility consequence, that every refresh must fit installed power and cooling. The consequence for AI sustainability is equally direct: without Dennard scaling, each process node no longer delivers proportional power savings, which pushes efficiency work toward specialized accelerators, such as GPUs and the Tensor Processing Unit (TPU), that exploit architectural parallelism rather than transistor physics alone.

15.3.1 Power envelopes and the energy wall

A short training burst can trip substation breakers or overwhelm chillers even when its cumulative megawatt-hour budget is modest, whereas a month-long job on a smaller cluster may consume far more total energy while staying safely inside its power envelope. The energy account divides work by realized work per joule and applies facility overhead at the boundary. Physical infrastructure, however, imposes a separate instantaneous power constraint: the grid connection, electrical distribution, and cooling system limit how much active fleet can operate safely at once.

The two limits obey different arithmetic (figure 15.2). The power limit is instantaneous. A facility cannot draw more than its grid feed ($P_{\text{facility}} \leq P_{\text{grid}}$). Because IT power is facility power divided by PUE, that feed limit also caps the number of servers active at once ($N \leq P_{\text{IT}}/P_{\text{server}}$) regardless of how long a job runs. The energy account is cumulative, integrating average power over wall-clock time to give run energy ($E_{\text{total}} = P_{\text{avg}} \times T_{\text{job}}$). Exceeding electrical or thermal capacity can force power limiting, throttling, or protective shutdown. Cumulative energy instead enters the electricity and emissions accounts, which also depend on tariffs and the applicable grid emissions factors.

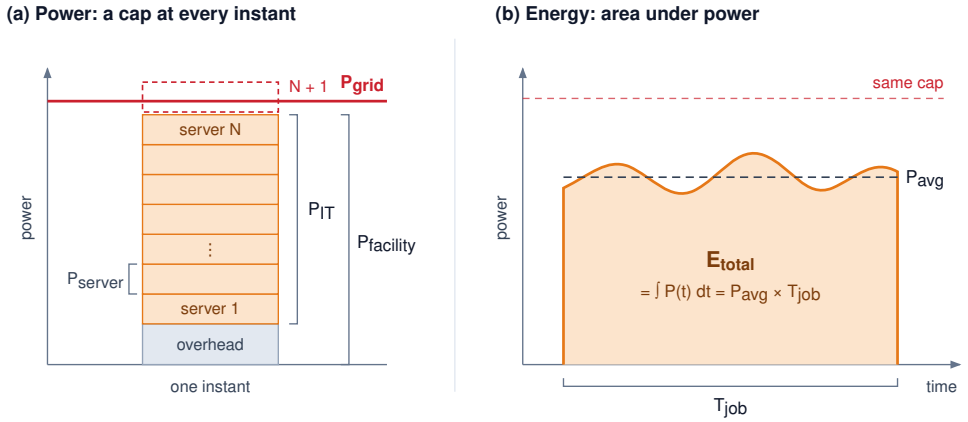

Figure 15.2: Power and Energy Boundaries: Substation grid feeds and facility cooling capacities impose instantaneous power caps on active accelerator concurrency. Integrating instantaneous power over wall-clock duration yields cumulative run energy; instantaneous power constraints and cumulative energy budgets represent distinct physical limits.

The energy wall is therefore a workload-specific feasibility problem, not a universal multiplier. Engineers must fit the required work, deadline, and quality target within available power and cooling, then calculate energy over the actual runtime. More efficient hardware helps only to the extent that the workload uses its efficiency and the resulting demand remains within those constraints. At frontier scale, the binding question is whether enough power can be delivered to the cluster at all.

> Lighthouse 15.1: Archetype A (GPT-4/Llama-3): The energy wall
>
> Training at the scale of Archetype A (GPT-4/Llama-3) exposes the power problem most sharply. At frontier scale, a single 25,000-GPU cluster is roughly 24 times the 1,024-GPU reference cluster used throughout this book. At 700 W per chip, it requires 17.5 MW of accelerator power before server, network, and cooling overhead. The constraint is physical, not financial: it is a grid capacity problem. Organizations operating Archetype A workloads may need dedicated power infrastructure or regions with excess renewable energy. Carbon-aware scheduling (shifting nonurgent work toward cleaner grid hours) and geographic optimization (placing workloads where power, cooling, and carbon constraints fit the job) therefore become fleet decisions as consequential as capacity planning.

That accelerator draw alone, before server, network, and cooling overhead, makes grid capacity a condition of deployment. Even when the grid can supply that sustained average power, the synchronized execution of a training job swings the load up and down across the cluster in unison.

15.3.2 Data center grid dynamics

When thousands of accelerators enter a collective communication barrier simultaneously and then resume matrix multiplication together, their power draw steps down and up in lockstep across the data hall. Unlike conventional cloud workloads whose independent requests smooth out across racks, synchronized training phases correlate power swings and push those transients toward the electrical grid interface.

War Story 15.1: When the grid became the bottleneck (2022)

Context: In July 2022, the Government of Ireland published its *Government Statement on the Role of Data Centres in Ireland's Enterprise Strategy*, addressing electricity constraints on further development in a country where data centers had become a substantial source of demand (Government of Ireland 2022; International Energy Agency 2024a).

Mechanism: Demand growth outpaced grid expansion. The statement reported that Dublin needed significant transmission reinforcement before it could accommodate new data-center power requests.

Impact: Grid connection capacity became the gating input for proposed Dublin facilities. New requests had to wait for transmission reinforcement or look to other regions, while contracted projects and existing sites continued to operate.

Fix: The statement prioritized network reinforcement, suitable regional locations, and development aligned with additional renewable supply. These were proposed responses to the constraint, not evidence that operators had already resolved it.

Systems lesson: Sustainable AI is constrained by interconnection, geography, and power-system adequacy. Efficient accelerators help, but the fleet cannot scale faster than the grid that feeds it, moving the binding constraint from the chip to the substation.

Training phases can change accelerator demand rapidly, but a device-side load step does not propagate unchanged to the utility connection. Its amplitude depends on per-device power, the number of coordinated workers, and computation-communication overlap. Local regulators, decoupling capacitors, power converters, and uninterruptible power supply (UPS) controls shape the upstream response. **Energy buffering** supplies or absorbs energy during a load change. A buffer discharges when the cluster steps up into computation and recharges when it falls back into communication, so the utility connection sees a smaller swing than the accelerators draw. Its location, capacity, and control response must match the electrical boundary being protected.

↩ Transient filtering by voltage regulators, decoupling capacitors, and UPS controls is established in section 2.6.1.1.

15.3.3 Training-scale energy concentration

The grid and siting constraints become severe because large training campaigns concentrate energy demand into long, synchronized runs. The GPT-3 anchor established the total scale; the new systems concern is its temporal concentration and the additional energy from distributed-training communication.[6]

Concentration in time is what separates a training campaign from the same energy spread across a year of ordinary load. A frontier run holds the 17.5 MW accelerator draw of the Archetype A cluster, before servers, network, and cooling, for the weeks or months the run lasts. Its synchronized computation and communication phases swing that draw together across the cluster (section 15.3.2). The facility must therefore provision for the sustained peak rather than for the annual average, and the grid sees one large, correlated load rather than many small independent ones. Large generative models intensify the problem when successive generations increase parameter counts, token budgets, or both.

Over the decade after 2016, those successive generations pushed single-run pretraining energy past the daily output of a utility-scale power plant (figure 15.3). Across 122 standalone runs that

6 | **Training communication overhead**: Distributed training adds network transfers, synchronization stalls, and checkpoint I/O to the energy account. The fraction is workload- and topology-dependent: parallelism strategy, tensor sizes, link technology, utilization loss, and any extra iterations determine whether communication is a secondary term or the binding energy cost. An energy estimate must therefore meter or model these terms for the actual run rather than apply a universal percentage.

document both cluster hardware and training duration, Epoch AI models peak facility power from accelerator count, chip thermal design power (TDP), a 1.82× server-chassis allowance (the ratio of peak server power to accelerator TDP), and facility PUE, then integrates that draw over wall-clock runtime (Epoch AI 2026c; Cottier et al. 2024). Along the 23 frontier pretraining runs in that set, upper-bound energy grew about 3.0× per year, slower than the 4.6× compute trend because hardware advances raised modeled FLOPs per joule by roughly 1.5× per year. Even with those efficiency gains, frontier runs crossed 24,000 MWh, one full day of continuous output from a 1 GW power plant at rated capacity, and reached roughly 240,000 MWh for Grok 3 before commercial nondisclosure began withholding duration data on later releases.

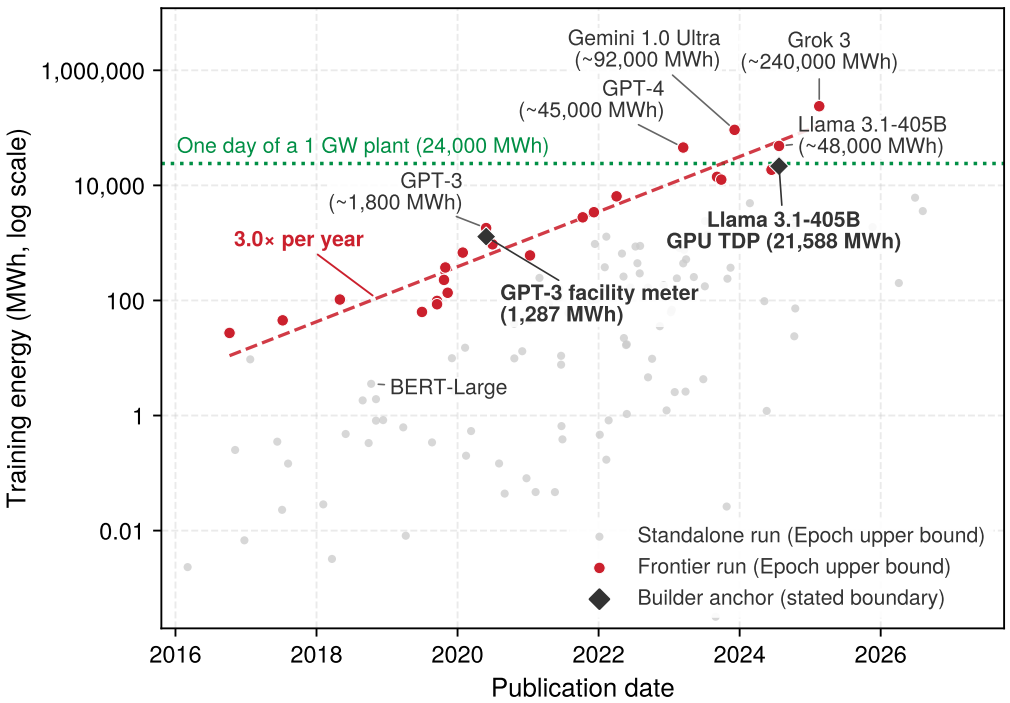

Figure 15.3: Training Energy of Notable Models: Single-run pretraining energy across 122 standalone runs since 2016, on a logarithmic axis. Gray dots and red circles plot peak-capacity facility upper bounds; the 23 frontier runs in red grow about 3.0× per year along the dashed least-squares fit (95 percent confidence interval 2.5× to 3.6×), before commercial nondisclosure withholds duration data on most releases after early 2025. The green dotted line marks 24,000 MWh, one full day of continuous output from a 1 GW power plant at rated capacity. Dark diamonds place builder figures below their paired facility upper bounds at each source's stated boundary: facility-metered energy for GPT-3 (1,287 MWh at 1.10 PUE) and accelerator-TDP energy implied by reported GPU-hours before host and cooling overhead for Llama 3.1-405B (21,588 MWh). Upper-bound series from Epoch AI (2026c), licensed under CC BY 4.0, using the facility-power model of Cottier et al. (2024); builder figures from Patterson et al. (2021), Meta AI (2024), and Dubey et al. (2024).

Placing builder-reported anchors alongside those modeled upper bounds shows why an energy claim is incomplete without its measurement boundary. For GPT-3, Patterson et al. (2021) report 1,287 MWh at the facility meter from OpenAI's 330 W per-GPU server average and a 1.10 PUE, whereas Epoch AI's peak-capacity model yields roughly 1,800 MWh (1.4× higher) because its 1.82× server-chassis allowance exceeds metered host and network draw. On TPU clusters where communication and I/O also pull sustained chip draw well below TDP, peak-capacity estimates exceed metered energy by two to three times (Patterson et al. 2021; Cottier et al. 2024).

Boundary scope widens the gap further for Llama 3.1-405B, whose model card (Meta AI 2024) reports 30.84M H100 GPU-hours at 700 W chip TDP (Dubey et al. 2024); multiplying the two implies 21,588 MWh at the accelerator-TDP boundary before host CPUs, network fabric, and cooling, or roughly 24,000 MWh once a 1.10 PUE is included. Epoch AI's full-facility upper bound reaches roughly 48,000 MWh, or 2.2× the accelerator-TDP product, and that ratio is the product of two factors.

A 1.97× server-and-PUE multiplier lifts the accelerator-TDP product for Meta's reported GPU-hours to roughly 43,000 MWh, and a 1.14× longer inferred wall-clock window (35.1M GPU-hours) lifts that figure to the Epoch bound. Even at the narrower accelerator-plus-PUE boundary, a single frontier training run now concentrates roughly a rated gigawatt-day of electricity into a few months of continuous operation.

The scaling laws introduced in section 1.3 explain why that concentration keeps growing. Within the ranges studied by Kaplan et al. (2020), test loss falls along smooth power laws in compute, dataset tokens, and parameters (figure 15.4), and for an energy account the small compute exponent is the quantity of interest. It means each fixed fractional reduction in loss costs a multiplicative increase in the resource. Because required compute scales as the ratio of the old loss to the new loss raised to the power 20, lowering compute-limited loss by 10 percent takes 8.2× the compute, and the next 10 percent costs the same multiple again. The diminishing return in principle 9 therefore arrives at the facility as a multiplicative increase in training energy unless realized FLOPs per joule improve at the same pace. Within these studied regimes, quality gains were purchasable with hardware and data, which is why compute demand has outrun hardware efficiency. Beyond training, high-volume deployed systems such as large-scale recommender systems and generative services require continuous inference at scale, consuming energy even after training completes. The cumulative energy burden therefore depends on both the one-time training run and the sustained query volume that follows deployment.

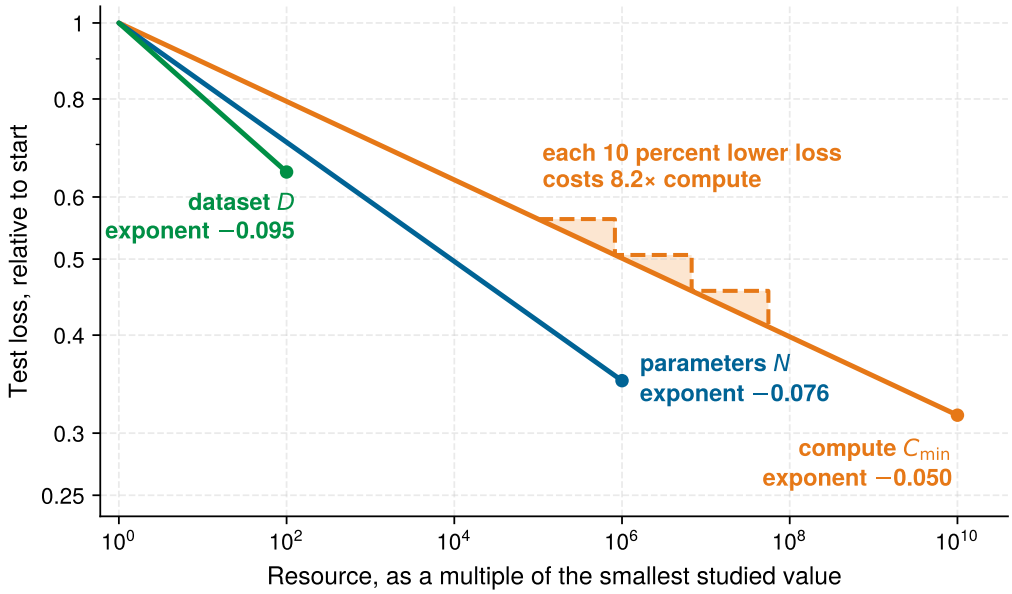

Figure 15.4: Model Scaling Laws: In the ranges studied by Kaplan et al. (2020), increasing model size, dataset size, and compute consistently reduced test loss. Each curve is normalized to its loss at the smallest studied value, so on these log-log axes its slope is its exponent; the dashed staircase on the compute curve marks the compute multiple that buys each successive fixed fractional reduction in loss. The figure reproduces source-paper notation, where N denotes parameter count and C_{min} denotes compute; elsewhere this book uses P for parameter count and O for operation count. These scaling laws explain why larger models, more data, and increased compute became attractive design choices within that regime.

Building out the clusters required by this concentration adds emissions before an operator runs a single training step, so expanding the fleet increases manufacturing impact even when each new accelerator uses less energy per useful task. An operator's inventory records that fabrication under Scope 3, the category for supply-chain emissions rather than on-site generation or purchased electricity (section 15.5.5 quantifies the manufacturing term, and section 15.5.2 assigns it to reporting scopes). Bounding any of these operational or embodied costs requires decomposing facility power into component-level models and direct telemetry, with every estimate declaring what it counts so that two estimates can be compared.

15.4 Energy Measurement and Modeling

A cluster drawing five megawatts during a large language model training run directs only a fraction of that power into matrix multiplications; host processors, memory, networking, power conversion, and cooling consume the rest. Engineers cannot optimize those losses from a single monthly utility bill.

Power delivery networks and cooling systems establish hard physical constraints on data center compute capacity. Systematic measurement turns these constraints into numbers an engineer can act on, provided the measurement fixes what it records and which lifecycle phase it targets before any power model or instrument is applied. Just as performance engineering requires profiling before optimization, sustainable AI engineering requires measurement before mitigation.

The decision procedure is the same throughout the chapter: measure the dominant lifecycle term, identify the physical bottleneck behind it, choose the intervention that changes that term, and check whether rebound effects erase the gain. Operational electricity, embodied carbon, cooling overhead, water use, and e-waste each call for different levers. Moving a job changes its operational setting without undoing emissions already incurred in manufacturing. Keeping hardware longer changes future procurement needs but also extends its operating energy account. Sustainable design therefore traces each intervention through the terms it changes before comparing the full result.

↩ The power-delivery path, thermal design power, and cooling limits of accelerators are established in section 2.6.1.

A sustainability claim is comparable only when it satisfies a **sustainability measurement contract**, which extends the reproducible measurement contract of section 9.3.1. That contract already fixes the workload and its quality gate, the system state, the timing boundary, and a statistical report with its window and spread. A sustainability claim adds five fields that fix what the energy figure covers and how it becomes carbon. The functional unit names the work being charged, such as one completed training run or 1,000 served requests at a specified output quality. The system boundary states which compute, memory, network, facility, manufacturing, and end-of-life terms enter that account (International Organization for Standardization 2006a, 2006b). Because multiple workloads and hardware generations share a facility, an attribution rule must then allocate shared infrastructure and embodied emissions. A declared average or marginal carbon signal converts the resulting energy into emissions. The time and location of operation fix which hours and which grid that signal describes. The statistical report carried over from the base contract must then also state how sensitive the conclusion is to those assumptions. A result that omits any of these fields may still diagnose one component, but it cannot support an end-to-end comparison. Large models released after GPT-3 often lack comparable energy disclosures, so filling these fields requires workload telemetry rather than published model documentation alone.

This contract separates two questions that otherwise become conflated. Attributional accounting allocates an inventory of observed emissions to workloads for reporting, typically using average grid intensity and the attribution rule above to share common infrastructure. Consequential accounting estimates how a decision changes emissions, for example by modeling the generator affected when a scheduler moves one megawatt-hour of demand. That generator is the plant on the margin of the grid's dispatch order, the plant whose output rises or falls when demand changes. A marginal emissions factor reflects that plant, which can burn fossil fuel even on a grid whose average mix is clean. Attributional accounting answers what share of an existing footprint belongs to a workload; consequential accounting answers whether changing the workload changes the footprint (Present et al. 2024). Carbon-aware scheduling therefore needs a signal appropriate for consequential analysis even when the organization's retrospective inventory uses attributional averages. The available marginal factors may be historical estimates, forecasts, or unavailable at the workload's required time and location. The contract must therefore record the source, resolution, and uncertainty rather than treat "marginal" as a universally observable datum.

15.4.1 Power modeling fundamentals

Lowering numerical precision or pruning half the weights of a model does not automatically cut the data center's electricity bill in half, because static leakage, memory traffic, and facility cooling each take their own share before power reaches the arithmetic units. Tracing where a watt goes requires a model that climbs three levels. The CMOS power equation at the chip level explains why voltage, precision, and switching activity matter. Physical optimization mechanisms show how algorithms

change those variables. Facility-level metrics reveal whether chip-level savings survive cooling and power-delivery overhead.

15.4.1.1 The CMOS power equation

Every digital circuit consumes power through two fundamental mechanisms. **Dynamic power** arises from switching transistors between states, while **static power** results from leakage current that flows even when transistors are nominally off. Equation 15.1 formalizes the total power consumption:

$$P_{\text{total}} = P_{\text{dynamic}} + P_{\text{leakage}} = \alpha C V^2 f + V I_{\text{leak}} \tag{15.1}$$

In the dynamic component $P_{\text{dynamic}} = \alpha C V^2 f$, the **switching activity factor** α summarizes capacitance-weighted switching per clock cycle and is not useful compute utilization. The term C is the effective load capacitance of the circuits and interconnect being switched. Supply voltage V enters quadratically, and clock frequency f sets the rate of switching opportunities. Useful work per joule depends on how much productive computation those transitions accomplish.

The static power component $P_{\text{leakage}} = V \cdot I_{\text{leak}}$ depends on leakage current, whose temperature sensitivity varies with device technology, voltage, and operating point. Rising temperature can increase leakage, adding heat that further raises temperature. Cooling and power control must keep this feedback within a stable operating range. A leakage model therefore needs parameters appropriate to the device rather than a universal temperature-doubling rule (Dayarathna et al. 2016).

The quadratic voltage dependence means that reducing voltage from 1 V to 0.8 V decreases dynamic power by 36 percent when activity, capacitance, and frequency stay fixed. The circuit must still complete its transitions within each clock period. A lower voltage may therefore require a lower frequency, which further reduces dynamic power but can lengthen execution. The energy benefit depends on both the resulting runtime and the static power consumed during it.

15.4.1.2 Physical mechanisms of energy optimization

The power equation illuminates why specific optimization techniques can improve efficiency. Quantization reduces numerical precision, which can shrink arithmetic units and reduce bytes moved through the memory hierarchy. The realized energy change depends on the implemented kernels, packing and conversion overhead, memory traffic, utilization, and the original bottleneck (Gholami et al. 2022). Lower precision does not by itself permit a lower supply voltage: voltage and frequency remain constrained by the physical implementation, process, and timing margins. Quantization therefore changes the operation and data-movement terms, but only measurement establishes the end-to-end reduction.

Pruning removes weights from neural networks, but it does not remove physical transistors from an existing accelerator. It saves dynamic energy when the resulting execution skips arithmetic or data movement. Structured pruning can map to smaller dense kernels; unstructured pruning needs an implementation that exploits the sparse pattern without spending the savings on indexing and irregular memory access. The power equation explains the benefit through reduced switching activity and powered execution time, not a change to the fabricated circuit capacitance C.

Specialized accelerators can improve useful work per joule by reducing instruction-control overhead, reusing data locally, and mapping operations to suitable datapaths. Keeping arithmetic units productive can amortize static and support power, but raising switching activity without raising useful output only raises power.

15.4.1.3 Facility-level power metrics

The PUE ratio has entered every facility-energy estimate since the opening energy boundary. The chip-level model now explains what happens inside the IT term, so the ratio itself needs its precise definition and its blind spots. Equation 15.2 gives the ratio at this chapter's facility boundary:

$$\text{PUE} = \frac{P_{\text{facility}}}{P_{\text{IT}}} \tag{15.2}$$

Definition 15.2: Power usage effectiveness (PUE)

Power usage effectiveness (PUE) is the data-center efficiency ratio ML system operators use to compare total facility power consumption against the power consumed specifically by IT equipment ($P_{\text{facility}}/P_{\text{IT}}$).

1. **Significance**: It measures the infrastructure overhead of the data center. A PUE of 1.0 is the theoretical ideal; a PUE of 1.10 means that for every 100 watts consumed by IT equipment, an additional 10 watts are required for non-IT facility overhead, including cooling and facility-side power distribution.
2. **Distinction**: Unlike computing efficiency (which focuses on FLOP/s per watt), PUE focuses on facility efficiency: it measures non-IT overhead relative to IT consumption. Idle equipment and power losses inside the IT measurement boundary remain in the denominator; PUE does not measure useful computation.
3. **Common pitfall**: A frequent misconception is that a low PUE means a "green" data center. PUE measures only overhead efficiency, not the carbon intensity of the energy source; a coal-powered data center can have a better PUE than a solar-powered one while having a much higher environmental impact.

Napkin Math 15.3: PUE: The cost of cooling

Problem: A team operates a 2 MW cluster. If the facility can be optimized from a legacy air-cooled PUE (1.58) to state-of-the-art (1.10), how much energy and money does that save annually?

Math: Energy saved is the difference in infrastructure overhead $(\text{PUE} - 1)$ across the IT load.

1. **Overhead Reduction**: 1.58 - 1.10 = 0.48.
2. **Annual energy savings**: 2 MW $\times$ 0.48 $\times$ 8760 h/year $\approx$ 8,409.6 MWh.
3. **Financial Savings**: 8,409.6 MWh $\times$ \$70/MWh $\approx$ \$588,672. This facility-only saving leaves the model unchanged and scales with reducible non-IT overhead, so cooling optimization matters most in inefficient facilities.

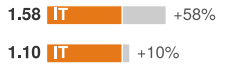

PUE measures facility overhead on top of IT energy.

A PUE of 1.0 would mean no measured non-IT overhead, not that all electricity produces useful computation. Idle servers and stalled accelerators remain inside the IT denominator. Surveyed data centers reported an average PUE near 1.55 in 2022, meaning that roughly 55 percent additional energy beyond IT consumption went to infrastructure, with older air-cooled sites higher still (Uptime Institute 2022). Leading hyperscale facilities achieve PUE between 1.1 and 1.2 through advanced cooling techniques including free-air cooling in cold climates, liquid cooling for high-density GPU clusters, and optimized power distribution.

Equation 15.3 formalizes **water usage effectiveness (WUE)**, capturing the water consumption that evaporative cooling and other processes require:

$$\text{WUE} = \frac{W_{\text{water}}}{E_{\text{IT}}} \tag{15.3}$$

The units are liters per kilowatt-hour, with typical values ranging from 0.5 to 2.0 L/kWh depending on climate and cooling technology. A data center with WUE of 1.8 L/kWh training a model requiring 10,000 MWh of IT energy would consume 18 million liters of water. That is equivalent to roughly 40–47 US household-years of water use under a 380,000–450,000 L/year household baseline. A facility-energy estimate must first be converted to the IT boundary before applying this WUE.

Serving has a water account too, and it is computed the same way. Google reports a fleet WUE of 1.15 L/kWh for the data centers that serve its Gemini assistant (Elsworth et al. 2025). The same study measured the energy of the median Gemini text prompt across the serving stack (section 15.7.3 examines that measurement). Applied to that prompt's IT energy, which excludes facility overhead,

the WUE gives the 0.26 mL of water per prompt that the study reports. A single prompt's water is negligible. A service answering billions of prompts a day turns the same arithmetic into a facility-scale demand, which is why the cooling choice, and the watershed it draws on, belongs in the serving design rather than only in the training budget. How a facility operates the cooling plant it already has also changes these facility-level metrics.

15.4.1.4 DeepMind energy efficiency case study

Google's cooling systems illustrate how operating settings affect PUE (2023). Server load, weather, and interactions among pumps and chillers complicate fixed operating heuristics (Buyya et al. 2010). Building on earlier work predicting PUE (Gao 2014), DeepMind trained neural networks from facility telemetry to predict future PUE and recommend settings. Additional models predicted temperature and pressure to check those recommendations against operating constraints.

In the reported live-site test, the system reduced cooling energy consumption by 40 percent and overall PUE overhead by 15 percent (Barroso et al. 2019; Evans and Gao 2016). By equation 15.2, a 15 percent cut in overhead removes 15 percent of the non-IT term (PUE minus 1), not 15 percent of total facility energy. At equal IT energy, the same fractional overhead reduction saves more electricity where the initial overhead is larger. The reported result establishes a benefit for the tested facility, not an identical saving at every site.

The saving came from changing how the existing pumps, chillers, and cooling towers were operated rather than from installing new cooling equipment, which makes control software a lever available to a facility that cannot be rebuilt. The transferable lesson is to evaluate operating controls alongside equipment changes. Telemetry supports the prediction, operating limits constrain the recommendation, and metered energy establishes whether the change saves electricity.

15.4.2 Systematic energy metrics

Cutting a model's FLOP count in half saves almost no energy when the workload spends nearly all of its joules shuttling weights between off-chip DRAM and the processor. Comparing hardware architectures and algorithmic changes therefore requires a systematic metric ladder that moves from energy per operation to energy per byte and then to the roofline relationship between arithmetic intensity and data movement.

15.4.2.1 Energy per operation

A basic metric for computational energy efficiency is energy consumed per operation, typically measured in picojoules. State the operation type and precision: a floating-point multiply-accumulate (MAC) is conventionally two FLOPs, whereas an integer MAC is two integer arithmetic operations.

Hardware architecture changes which overheads can be reduced. Each architecture class carries an effective energy coefficient in this chapter's models, paired with the measurement needed to test it (table 15.1). The coefficients span three orders of magnitude, from 100 pJ/FLOP for a general-purpose core, where instruction fetch and control dominate each operation, to 0.1 pJ/FLOP for a fixed-function datapath. They fold local operand delivery and control into each operation, which is why the GPU coefficient of 10 pJ/FLOP sits well above the roughly 3.7 pJ of a bare FP32 multiply at 45 nm (Horowitz 2014). They are order-of-magnitude modeling inputs, not a ranking of products.

Table 15.1: Architecture and Energy Measurement: The effective-energy column gives the order-of-magnitude coefficient this chapter's energy models assign to each architecture class, including local operand delivery and control. Specialization can reduce execution overhead, but device specifications with different operation types and system boundaries do not establish an end-to-end energy ranking. Compare the same useful task at the required quality.

Architecture	Effective energy (model)	Potential energy advantage	Comparison must include
General-purpose CPU	100 pJ/FLOP	Flexible execution for diverse workloads	Instruction overhead and achieved utilization
GPU	10 pJ/FLOP	Parallel arithmetic and data reuse	Precision, occupancy, and memory traffic
Systolic accelerator	1 pJ/FLOP	Local reuse across a regular datapath	Supported shapes, precision, and data movement

Architecture	Effective energy (model)	Potential energy advantage	Comparison must include
Edge accelerator or fixed-function ASIC	0.1 pJ/FLOP	Specialized low-power execution	Host, memory, peripherals, and duty cycle
Analog or in-memory compute	Not modeled	Reduced movement for supported operations	Conversion, peripheral circuits, and achieved quality

Precision affects energy per operation through implementation-specific arithmetic and data movement. Narrower formats can reduce circuit area, switched capacitance, and memory traffic, but bit width alone does not determine wall energy. The measurement contract requires a matched workload and quality target, with the same memory and supporting-system boundary.

15.4.2.2 Energy per byte

In AI workloads with low arithmetic intensity, moving operands across storage and interconnect boundaries consumes far more energy than the arithmetic units themselves. An illustrative cost model across registers, caches, main memory, solid-state storage, and the network makes this hierarchy concrete; its coefficients teach the relative cost of movement across boundaries rather than fixed measurements of every platform.

With the modeled coefficients of table 15.2, movement from DRAM can dominate the arithmetic it feeds. The rows trace the path from registers through L1 cache and L2 cache to DRAM, Non-Volatile Memory Express (NVMe), and network transfers. Under the GPU coefficient of 10 pJ/FLOP from table 15.1, accessing one FP32 operand from DRAM (4 bytes × 160 pJ/byte = 640 pJ) costs 64× the energy of the operation that consumes it.

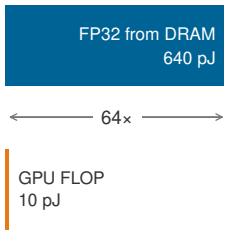

Fetching an FP32 operand from DRAM costs 64× the FLOP that uses it.

These coefficients agree with the movement costs priced per bit in section 2.3.1.1 once the units match. That section's off-package DRAM figure of about 20 pJ per bit is this table's 160 pJ/byte, so one 16-bit operand fetched from DRAM costs 320 pJ, where this example fetches a 4-byte FP32 operand. The bare FP16 multiply-accumulate of about 1.5 pJ quoted there sits below the effective GPU coefficient here because that coefficient folds control and local operand delivery into each operation. Transfers across the network fabric are priced per bit in section 5.2, whose fabric-path figure of about 35 pJ/bit (280 pJ/byte) charges only the serializers, transceivers, and switches between endpoints. The network row here is a whole-system coefficient that also charges endpoint memory, host software, and idle network power, so the two differ by more than an order of magnitude.

Table 15.2: Illustrative Movement-Energy Coefficients: The model assigns energy per byte at each stated boundary. Real values depend on transfer size, locality, link or device implementation, and supporting-system coverage.

Memory level	Energy cost (pJ/byte)
Register	0.1 pJ/byte
L1 Cache	1 pJ/byte
L2 Cache	5 pJ/byte
DRAM	160 pJ/byte
NVMe SSD	1000 pJ/byte
Network	> 10000 pJ/byte

The resulting design levers target data movement as well as arithmetic. On-chip SRAM and registers support operand reuse, while systolic arrays pass operands and partial sums between neighboring processing elements instead of repeatedly returning to off-chip memory. Compression and sparse representations can reduce transferred data when saved payload exceeds metadata and conversion overhead. The energy benefit therefore depends on the implemented memory traffic, rather than solely on nominal model size or arithmetic throughput.

15.4.2.3 Arithmetic intensity and energy roofline

Arithmetic intensity, the operations-per-byte ratio recalled in section A.3, connects the operation count to the data-movement volume. Combined with energy per operation and per byte, it identifies

↩ The performance roofline and its ridge point are established in section 2.3.2.

which contribution dominates this energy model; it does not by itself identify the resource limiting execution speed. Equation 15.4 defines the workload ratio. In these equations, O is the operation count in FLOPs, D_{vol} is data volume in bytes, E_{compute} is energy per FLOP, and E_{move} is energy per byte moved:

$$I = \frac{O}{D_{\text{vol}}} \tag{15.4}$$

At fixed energy coefficients, arithmetic intensity measured in FLOP/byte determines the relative size of the compute and movement terms. Equation 15.5 expresses modeled energy E_{total} at the chip-and-memory boundary as their sum (the facility total E_{facility} appears later in equation 15.9), while equation 15.6 isolates the roofline-style dominant term. Static power and supporting systems require additional terms if the comparison includes them:

$$E_{\text{total}} = O \times E_{\text{compute}} + D_{\text{vol}} \times E_{\text{move}} \tag{15.5}$$

$$E_{\text{dominant}} = \max\left(O \times E_{\text{compute}}, D_{\text{vol}} \times E_{\text{move}}\right) \tag{15.6}$$

The maximum identifies the larger energy contribution, not the sum. Near balance, omitting the other term substantially understates the modeled energy. Equation 15.7 defines the crossover arithmetic intensity where compute and memory energy balance:

$$I_{\text{crossover}} = \frac{E_{\text{move}}}{E_{\text{compute}}} \tag{15.7}$$

For a GPU with E_{compute} of 10 pJ/FLOP and E_{move} of 160 pJ/byte (DRAM access):

$$I_{\text{crossover}} = \frac{160 \text{ pJ/byte}}{10 \text{ pJ/FLOP}} = 16 \text{ FLOP/byte}$$

The **energy roofline model** in figure 15.5 bounds energy efficiency (GFLOP/J) against arithmetic intensity (FLOP/byte), with a 16 FLOP/byte crossover under the stated component costs. Below this point, the larger energy contribution is memory movement (160 pJ/byte). Above it, arithmetic sets the horizontal upper bound of 100 GFLOP/J (10 pJ/FLOP). The piecewise bound keeps only the larger contribution, whereas actual energy includes both and any omitted system overhead. Its crossover therefore answers a different question from the performance ridge, which equals peak arithmetic throughput divided by memory bandwidth and often sits well above 16 FLOP/byte on modern accelerators. A kernel whose arithmetic intensity falls between the energy crossover and the performance ridge is still paced in execution time by memory bandwidth even though arithmetic already accounts for more of its modeled dynamic energy. Applying this crossover to matrix multiplication shows how problem dimension alone flips the dominant energy term.

Napkin Math 15.4: MatMul energy analysis

Problem: How do matrix dimensions N dictate whether arithmetic energy E_{compute} or memory energy E_{move} dominates matrix multiplication $\mathbf{C} = \mathbf{AB}$?

Math:

1. Operation & byte count: $2N^3$ FLOPs and $12N^2$ bytes at FP32, assuming ideal reuse: read each input once and write the output once at the modeled memory boundary. Tiling and limited capacity can increase measured traffic.
2. Arithmetic intensity: $I = \frac{N}{6}$ FLOP/byte, so the matrix sits exactly at the energy crossover when $N = 96$ ($I = 96/6 = 16$), the small-matrix row below.
3. Regime comparison: small dense work sits at the crossover, large dense work is dominated by arithmetic, and element-wise work by memory movement (table 15.3).

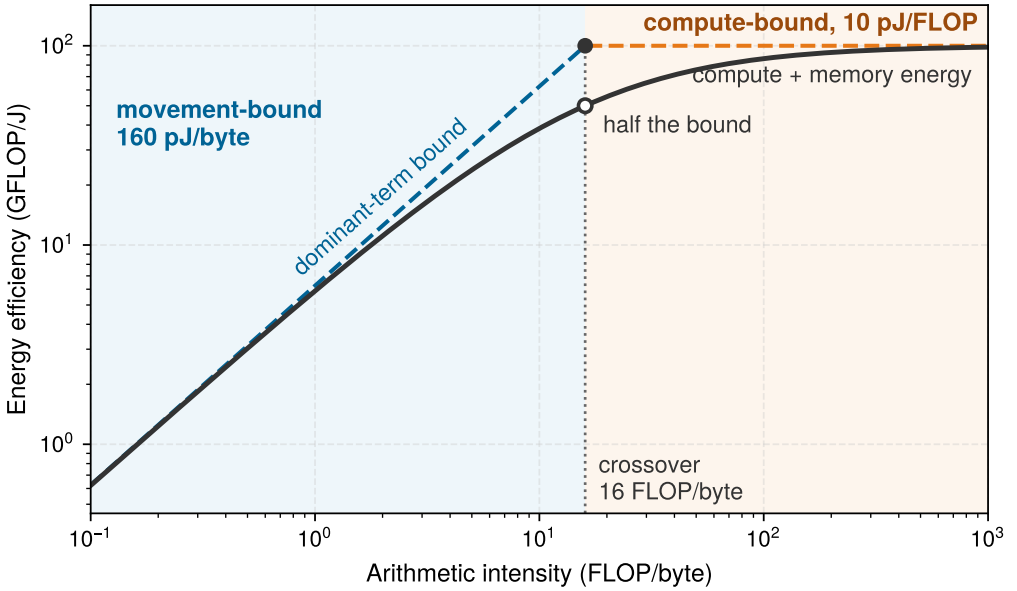

Figure 15.5: Energy Roofline Model: Just as performance rooflines limit FLOP/s based on bandwidth, energy rooflines limit FLOP/J. The dashed curve uses the dominant energy term (blue where data movement dominates, orange where arithmetic dominates); the solid curve includes both compute and memory energy using the same coefficients as the worked example. At the crossover both terms contribute equally, so total-energy efficiency is half the dominant-term bound. Static and other system energy require additional terms.

Table 15.3: Energy Bottleneck Comparison: The same energy model classifies small dense work, large dense work, and element-wise work by their dominant energy term.

Workload	Arithmetic intensity	Compute energy	Memory energy	Optimization priority
Small matrix ($N = 96$)	$I = 96/6 = 16$ FLOP/byte	$2 \times 96^3 \times 10$ pJ = 17.69 µJ = 0.0177 mJ	$3 \times 96^2 \times 4 \times 160$ pJ = 17.69 µJ = 0.0177 mJ	Balanced at the crossover
Large matrix ($N = 1000$)	$I = 1000/6 = 167$ FLOP/byte	$2 \times 10^9 \times 10$ pJ = 20 mJ	$3 \times 10^6 \times 4 \times 160$ pJ = 1.92 mJ	Improve compute efficiency
Vector addition ($N_{\text{vec}} = 1000$)	$I = 1000/12000 = 0.083$ FLOP/byte	1000×10 pJ = 0.00001 mJ	12000×160 pJ = 0.00192 mJ	Reduce data movement through fusion

Systems insight: Energy optimization follows the same bottleneck logic as latency optimization. Dense matrix multiply rewards efficient arithmetic; element-wise work rewards reducing memory movement.

The energy roofline model reveals why different optimization strategies suit different workloads. When arithmetic dominates the energy account, large dense matrix operations benefit from lower energy per executed operation. When data movement dominates, element-wise kernels benefit from data layout optimization, kernel fusion to reduce memory round-trips, and on-chip memory utilization. Faster execution alone does not establish either saving; average power and execution time must be measured together. This framework guides architectural and algorithmic choices for sustainable AI system design.

15.4.3 Energy measurement techniques

Reading `nvidia-smi` during a training step misses host DRAM, NVMe storage, network switches, and chiller pumps, whereas reading the substation utility meter lumps every co-located job into a single total. Quantifying AI system energy consumption therefore requires choosing measurement instruments across multiple levels of the hardware stack. On-chip counters provide fine attribution, mobile profilers expose platform-specific subsystems, external edge instruments capture duty-cycle behavior, and system-level tools connect component measurements to the facility overhead that device counters miss.

15.4.3.1 Hardware power counters

Modern processors expose manufacturer-provided power telemetry that software can query through device interfaces. These counters provide workload-correlated estimates within the rails and sampling windows the device reports; they are not universal ground truth for wall-plug energy, and their update rate, calibration, and component coverage vary by platform. External instruments remain necessary when a comparison requires traceable total energy or finer temporal resolution.

Intel's **running average power limit (RAPL)** interface exposes cumulative energy counters for supported power domains, with coverage depending on the processor. A controlled measurement samples the relevant counter before and after the workload and divides the energy delta by elapsed time to recover average power. This makes RAPL useful for host-side work, provided the account identifies the domains actually reported and handles counter rollover. Package and subdomain counters must not be added when their coverage overlaps. RAPL does not measure discrete GPU energy, and access can require elevated permissions.

NVIDIA GPUs expose power information through the NVIDIA Management Library, accessible via the `nvidia-smi` command-line tool or programmatic bindings. GPU power monitoring must distinguish an instantaneous reading from a sensor-averaged value; the reported behavior depends on the device and query. Integrating samples requires a cadence and workload interval appropriate to that averaging window. Sampling an averaged signal more frequently does not recover short power spikes. Supported accumulated-energy counters can simplify interval accounting, but their scope and resolution still need checking, and neither method supplies full facility energy by itself.

Edge devices and microcontrollers often lack built-in power counters and require external instrumentation. Select current range, voltage measurement, sampling rate, and analog bandwidth to capture the workload's active and sleep intervals. Board-level monitors can integrate inference energy; resolving shorter events requires adequate instrument bandwidth and careful treatment of filtering and aliasing. A fast sample rate alone cannot attribute energy to individual operations.

15.4.3.2 Mobile platform energy profiling

On mobile platforms, measurement depends on how much subsystem attribution the host operating system exposes. Some platforms report per-rail or per-component power for the CPU, GPU, neural processing unit, and radio, which lets a developer see which subsystem dominates total draw. Others expose only relative energy-impact scores or thermal states, which are useful for spotting regressions during development but cannot support an absolute energy claim. Vendor profilers that correlate power traces with kernel execution sit between these extremes. The design question is whether the reported quantity covers the electrical boundary the comparison needs, not which tool to use. When it cannot, external instrumentation is required. Edge power measurement instruments differ by measurement role, and each role requires its own checks (table 15.4).

Table 15.4: Edge Power Measurement Instruments: Choose instrumentation by the event duration, dynamic range, and electrical boundary being measured. Verify configured specifications rather than assuming one device is universally sufficient.

Instrument class	Measurement role	Required checks
Board current monitor	Integrate energy on selected rails	Voltage, sample interval, shunt losses, and event duration
Multi-rail monitor	Attribute energy among components	Channel synchronization and coverage

Instrument class	Measurement role	Required checks
Wide-range external analyzer	Capture sleep and active intervals	Bandwidth, range transitions, and calibration
Automated power analyzer	Repeat duty-cycle and battery tests	Triggering, integration window, and supply behavior

15.4.3.3 Edge measurement methodology

Edge energy measurements are useful only when they reflect the deployed duty cycle rather than an idealized active-only inference burst, so reproducible evaluation requires experimental controls matched to that workload. Simultaneous current and voltage logging must span both the active and quiescent states being measured. Logging voltage alongside current captures rail voltage fluctuations rather than inferring power from current alone. Thermal and frequency stabilization is necessary for sustained-load comparisons, whereas cold starts and periodic bursts must be measured across their actual state transitions. Because edge devices spend most of their deployment life waiting between events, realistic duty-cycle accounting records both peak inference draw and time-averaged power across typical idle intervals. Component attribution may isolate or subtract peripheral loads to expose accelerator consumption. A deployment estimate, however, must restore the radios, sensors, and other subsystems needed to deliver the service. For duty cycle accounting, equation 15.8 expresses this relationship between active and idle power:

$$P_{\text{avg}} = P_{\text{active}} \times \alpha_{\text{duty}} + P_{\text{idle}} \times (1 - \alpha_{\text{duty}}) \tag{15.8}$$

where α_{duty} is the duty cycle (fraction of time performing inference).

15.4.3.4 System-level energy profiling

Comprehensive energy accounting requires combining chip-level measurements with infrastructure overhead. The component terms must be disjoint: GPU-board telemetry that already includes device memory must not be added to another measurement of that same memory. They must also cover the IT boundary, with storage and server support loads included in the relevant measured terms. Equation 15.9 formalizes total energy as the sum of those component contributions scaled by facility overhead:

$$E_{\text{facility}} = (E_{\text{CPU}} + E_{\text{GPU}} + E_{\text{memory}} + E_{\text{network}}) \times \text{PUE} \tag{15.9}$$

No single counter spans the full energy path, so system-level profilers like Intel VTune, NVIDIA Nsight Systems, and open-source tools such as PowerJoular aggregate measurements across components. On-chip counters estimate device power and may differ in sampling rate, coverage, and calibration; a rack power distribution unit (PDU) provides an independent measurement of the downstream IT load. A rack PDU does not measure central chillers, pumps, lighting, or upstream conversion losses, so facility energy must instead come from facility meters or from an explicitly time-matched PUE applied to the measured IT load. The same rule holds for power: facility power is the measured IT power, the sum of server, network, and storage equipment power, multiplied by the PUE of equation 15.2, so the overhead enters once at the facility boundary.

Development-time carbon tracking matters when it feeds placement and complexity decisions rather than becoming a report after the fact. Tools such as CarbonTracker (Anthony et al. 2020) and CodeCarbon (Schmidt et al. 2021) wrap the workload boundary, estimate energy from hardware counters or utilization models, attach the local grid intensity, and record the resulting emissions alongside the experiment metadata. The important systems pattern is not the API call but the timing of the measurement: the estimate must appear while model size, training duration, region, and schedule are still adjustable, before the run has already consumed its energy. Even then, facility energy is still not a complete footprint, because converting it to emissions requires a grid factor and the hardware that drew it carries manufacturing emissions that no power meter records.

15.5 Carbon Footprint Calculation

Consider a data center supplied by a hydro-heavy grid. Its operational carbon intensity can be far below a fossil-heavy grid, but hydroelectric generation is not lifecycle-zero, and the AI system is not carbon-free. Reservoir construction and operation have impacts, while mining the silicon, manufacturing the GPUs, and pouring the concrete for the data center emitted carbon before the servers were turned on. A complete footprint calculation must account for operational electricity, embodied carbon, and the ecological effects excluded from a carbon-only metric.

Building that footprint term by term starts by asking which lifecycle phase dominates the workload and then whose inventory each emission belongs to, because the same physical emission can appear in more than one organization's report. The grid factor, operational carbon, embodied carbon, and their lifecycle sum then each take an equation that generalizes across workloads, grids, and amortization assumptions.

15.5.1 Carbon accountability and lifecycle phases

Published estimates of the energy of a single chat prompt vary by about an order of magnitude because different studies count different components of the serving stack (Elsworth et al. 2025), and best-effort estimates stripped of their assumptions quickly circulate as facts (Luccioni et al. 2025). Without a workload-level accounting model that fixes those boundaries, teams cannot audit whether choosing a smaller model, lower-power hardware, a cleaner region, or deferred training actually reduces environmental impact.

Carbon footprint analysis turns sustainability from a general obligation into an auditable design constraint by linking energy consumption, grid carbon intensity, and lifecycle resource demands to the same decisions that govern performance and efficiency (Schwartz et al. 2020; Patterson et al. 2021). Traceability is the technical bridge between measurement and accountability, and it is the engineering response to a transparency gap: the absence of per-model energy, carbon, and water records that would let anyone outside the team check a claim. A team should be able to connect a model version, training run, serving workload, region, energy source, and lifecycle estimate well enough that carbon and water claims can be verified.

For measurement to constrain design, reported metrics must expose workload-level cost rather than aggregate cloud-scale claims, which stay opaque when reporting holds at the company-total level. The production measurement in section 15.7.3 shows how far the boundary alone moves a reported figure. The practical standard is therefore accountability for resource usage across the full AI lifecycle. It requires the same evidence chain used for latency, accuracy, and reliability: a claim should connect to logs, meters, model versions, and deployment decisions rather than to aggregate sustainability pledges alone. The footprint calculation makes that evidence chain explicit by combining workload energy, facility PUE, grid carbon intensity, and embodied carbon, and it starts from the lifecycle phase that dominates.

15.5.1.1 Three-phase lifecycle assessment framework

The practical question is which lifecycle phase dominates a given workload, because the answer determines the optimization lever. Effective carbon footprint measurement therefore separates three phases that collectively determine environmental impact: training, inference, and the manufacturing of the hardware both run on. Two other groupings appear later in the chapter, the reporting scopes of section 15.5.2 and the lifecycle stages of section 15.8. Both regroup these same emissions, by reporting organization and by time, rather than adding new ones. Table 15.5 maps the three groupings onto one another.

For training-centric research workloads, the training phase often dominates operational emissions because mathematical optimization requires sustained parallel computation.[7] Geographic placement then changes the operational result, because an identical workload can encounter very different grid intensities.[8] For high-volume production services, by contrast, the inference phase can dominate lifetime emissions because serving repeats continuously after training ends, so its cumulative impact scales with deployment breadth and usage frequency. The deciding comparison is cumulative serving emissions over the deployment period against training emissions, including the experimental runs that preceded the final one. The more widely a model is served and the longer it stays deployed,

[7] **Optimizer memory as energy cost**: Adaptive Moment Estimation (Adam) stores per-parameter first and second moment estimates, unlike plain stochastic gradient descent (SGD). For a 70B model in FP32, these two moment buffers total about 560 GB, bringing parameters plus moments to about 840 GB (excluding gradients). The sustainability implication is direct: larger optimizer state means more off-chip memory traffic per training step, and at roughly 160 pJ/byte for a DRAM access, memory movement can dominate the energy budget of parameter updates.

[8] **Carbon intensity variance**: Regional averages and time-varying grid factors answer different questions. A generation technology's lifecycle factor is neither a measured regional average nor the marginal emissions response to another job. Scheduling must choose the signal appropriate to its decision and retain the source, period, and boundary.

the more likely inference is the larger term, and section 15.7.3 computes how many days of serving at a fixed query volume overtake a training run.

The manufacturing phase contributes embodied carbon from hardware production, including material extraction, semiconductor fabrication, and supply chain logistics.[9] For a fixed manufacturing footprint, longer operation on a carbon-intensive grid increases operational emissions and lowers the embodied share of the total. Cleaner electricity or less operating time raises that relative share without changing the manufacturing emissions already incurred. Operational efficiency does not undo those emissions; procurement, design, reuse, and service-life decisions determine future embodied costs.

15.5.2 Emission scopes and the development boundary

When an AI organization rents a cloud GPU cluster, the electricity that powers the run appears on the cloud provider's utility bill, the chip fabrication emissions appear in the foundry's inventory, and neither registers as direct fuel or electricity on the renter's own meters. Whereas the training, inference, and manufacturing account groups emissions by the physical activity that produces them, the three emission scopes defined by the **Greenhouse Gas (GHG) Protocol** (World Resources Institute and World Business Council for Sustainable Development 2023) classify that same footprint relative to a reporting organization. These are different views of the same emissions, not additional totals to add together. Connecting the workload account to an organizational inventory requires naming who owns the equipment, purchases the electricity, and buys the service.

Within an owned data-center facility boundary, Scope 1 emissions include direct combustion in owned generators and refrigerant leakage from cooling equipment, whereas purchased electricity used to run cooling or IT equipment belongs in Scope 2, not Scope 1. Scope 2 emissions represent indirect emissions from purchased electricity, steam, heat, and cooling consumed by the reporting organization. Purchased electricity is often the dominant category for an operator's owned facility, and it varies dramatically by geographic location and grid energy mix. Electricity used by a cloud provider is Scope 2 for that provider, but a customer's purchased cloud service is generally accounted for in the customer's Scope 3 value chain. The physical electrons do not change; the reporting boundary does. The geography notebook in section 15.2 uses a declared pair of illustrative factors; actual accounting requires factors matched to the location, time, and method.

Scope 3 emissions constitute the most complex category, encompassing hardware manufacturing, cloud supply chains, transportation, disposal, and downstream use. Semiconductor manufacturing is carbon-intensive.[10] The manufacturer reports the same fabrication emissions inside its own organizational boundary, so the two inventories describe one physical emission from two sides. How long operation takes to match that embodied term depends on how hard and on which grid the hardware runs. The H100 break-even worked with the lifecycle phases (section 15.5.1.1) stretches from a few weeks at rated power on the US-average grid to months or years at low utilization or on a clean grid. Under company-wide or value-chain AI accounting, Scope 3 can dominate even when Scope 2 dominates a single owned facility's operational footprint.

Downstream use can also enter Scope 3, including end-user smartphones, IoT devices, and edge computing.[11] Companywide totals do not by themselves identify the footprint of a particular AI service; that attribution still requires workload-level accounting.

The GHG Protocol arranges these scopes as concentric reporting boundaries that widen outward from the operator's direct control (figure 15.6). On-site diesel generators sit in Scope 1, purchased grid electricity for compute and cooling in Scope 2, and upstream semiconductor fabrication and end-of-life recycling in Scope 3.

The operational emissions of a final training run and its deployed service omit the development that produced them, including failed experiments, hyperparameter searches, and the build and test infrastructure. Repeated experiments can exceed the cost of the final run. A search that trains k candidates at a fraction f of the final run's cost adds kf final-run equivalents before the chosen model starts, so a sweep of 50 candidates at 10 percent of full scale costs 5× the run that gets reported. Continuous integration, compilation, and evaluation add further work whose contribution depends on run frequency and scale. The development boundary must record these activities before assigning their emissions to the reporting organization's inventory; the phase of model development does not itself determine the emissions scope.

[9] **Embodied carbon break-even**: Embodied carbon covers manufacturing, transport, and end-of-life treatment rather than use. Allocating NVIDIA's HGX H100 baseboard product carbon footprint evenly across its eight GPUs gives roughly 164 kg CO_2e per H100 (NVIDIA Corporation 2025); assuming continuous accelerator power of 700 W on the average U.S. grid, the modeled operational emissions match that allocation in roughly three to four weeks. This comparison omits host and facility overhead and uses rated power rather than a measured workload trace. A cleaner grid lengthens the modeled interval, but this equality alone does not decide whether replacing hardware reduces future emissions.

[10] **Lithography and embodied energy**: Advanced semiconductor fabrication can use energy-intensive lithography, including extreme ultraviolet (EUV) processes. Emissions per usable chip depend on the process sequence, facility energy mix, yield, die area, and allocation method. Neither transistor count nor a nominal process-node label alone determines the number of exposure steps or manufacturing footprint.

[11] **Edge AI energy trade-off**: Edge inference can reduce network transfers and response latency, while distributing computation across devices with different duty cycles and power states. It does not eliminate network or cloud work, and distributed consumption, although harder to attribute, remains part of the accounting boundary; section 15.7.4 works through the comparison.

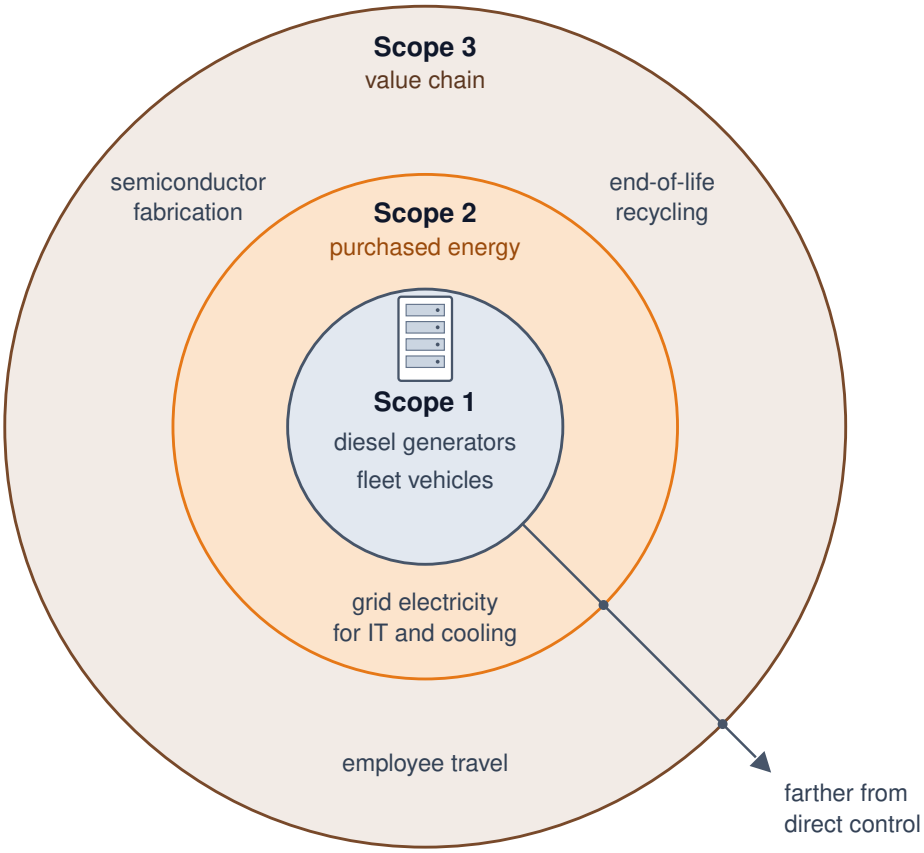

Figure 15.6: GHG Emission Scopes: Concentric rings place the reporting organization's data center at the center, with Scope 1 (direct emissions) innermost, Scope 2 (purchased energy) around it, and Scope 3 (value chain) outermost. Representative sources sit inside each ring. The rings show reporting distance from the operator's direct control, not relative size. Scope definitions follow the GHG Protocol (World Resources Institute and World Business Council for Sustainable Development 2023).

With the scopes and the development boundary defined, the activity phases of section 15.5.1.1, the lifecycle stages of section 15.8, and the scopes can be laid side by side (table 15.5). Each row is one set of emissions seen three ways, so a total built from one column must never be added to a total built from another.

Table 15.5: Three Views of One Footprint: The activity phases group emissions by the work that causes them, the lifecycle stages by when they occur, and the GHG scopes by which organization reports them. The same physical emissions appear once in each view; the scope assignment changes with the reporting organization, not with the hardware.

Activity phase	Lifecycle stage	Scope for an owner-operator	Scope for a cloud customer
Training (development runs and the final run)	Design and experimentation; use	Scope 2 for purchased electricity; Scope 1 for on-site generation	Scope 3 (purchased service)
Inference (serving)	Use	Scope 2 for purchased electricity; Scope 1 for on-site generation	Scope 3 (purchased service)
Manufacturing (embodied carbon)	Manufacture; disposal	Scope 3 (capital goods and end-of-life treatment)	Scope 3 (within the purchased service)

Applying the scope framework in practice requires classifying the emission sources across a typical ML platform's lifecycle.

Checkpoint 15.2: Accounting for invisible carbon

Audit the carbon footprint of a machine learning platform by classifying the following emission sources into Scope 1 (Direct), Scope 2 (Indirect Energy), or Scope 3 (Value Chain). State the reporting organization before assigning each scope:

- ☐ **Backup generation**: classify diesel burned by backup generators during a grid outage at an owned facility.
- ☐ **Leased cloud cluster**: classify the electricity for both reporting entities, the cloud operator that purchases it and the customer that purchases the compute service.
- ☐ **Embodied manufacturing**: classify the embodied carbon emitted during the manufacturing of the GPUs by TSMC.
- ☐ **End-user operation**: classify emissions from the end-user's smartphone battery while running a deployed mobile inference app.

15.5.3 Grid carbon intensity and regional variation

Two reports can price the exact same megawatt-hour of data-center electricity at carbon intensities that differ several-fold simply by choosing a different emissions factor or boundary. Generation-technology lifecycle factors include fuel supply, construction, operation, and decommissioning, whereas regional grid factors reflect an electricity mix over a stated place and period and may exclude upstream lifecycle emissions. Each type of factor must be identified before it enters the calculation (table 15.6).

Table 15.6: Electricity Emissions Factor Boundaries: A technology factor is not a regional grid measurement. Record the period, location, imports, greenhouse gases, and lifecycle coverage before comparing or multiplying a factor by workload electricity.

Factor	What it describes	Appropriate use
Generation lifecycle	Emissions across the lifecycle of an electricity technology	Comparing generation options under consistent lifecycle assumptions
Grid average	Emissions attributed to an electricity mix over a place and period	Attributing workload electricity emissions with a declared operational or lifecycle boundary
Grid marginal	Estimated change in emissions when demand changes	Evaluating a proposed shift or reduction in electricity demand

For its 2024 data centers, Google reports a location-based grid factor, the average intensity of the grids its data centers draw from, of 345 g CO_2e/kWh and a market-based factor of 94 g CO_2e/kWh after crediting its purchases of carbon-free electricity, a 3.7× difference for the same electricity (Elsworth et al. 2025). The median Gemini text prompt whose water use section 15.4.1.3 computed therefore carries about 0.08 g CO_2e or 0.023 g CO_2e of electricity emissions depending on the factor, and only the second matches the figure the company publishes. Neither is wrong, but a comparison is meaningful only when both sides declare the same factor type, period, and boundary.

Applied to a fixed 10,000 MWh facility-energy account, regional electricity factors change operational carbon by nearly two orders of magnitude (figure 15.7). The assumed factors range from 10 g/kWh for Norway (100 t CO_2 for the workload) to 820 g/kWh for Poland (8,200 t CO_2), producing a calculated 82× emissions ratio. These rounded annual-average factors (International Energy Agency 2023), the same ones notebook 15.2 used for Quebec and Poland, illustrate the effect of placement while holding facility energy constant. They are not measurements of the same job run at every site. Actual placement decisions require time-matched grid evidence and must include changes in facility overhead and the energy needed to move or resume the work.

15.5.4 Operational carbon calculation

Multiplying raw GPU thermal design power by runtime and a regional grid factor understates a job's operational carbon when it omits host servers and chiller power, or overstates it if facility overhead is accidentally applied a second time. Equation 15.10 computes operational carbon from

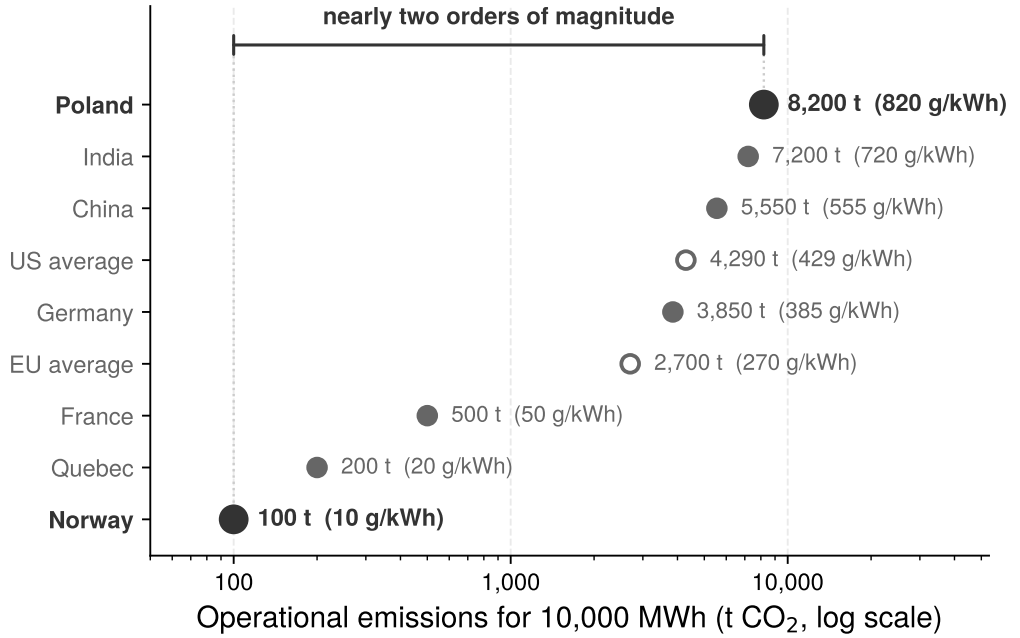

Figure 15.7: Comparing Regional Electricity Emissions: Operational carbon of one fixed facility-energy workload under rounded annual grid-average factors for each region (International Energy Agency 2023); the Quebec factor is the utility's reported mix. With facility electricity held fixed, the regional generation mix alone changes operational carbon by nearly two orders of magnitude. Open circles mark regional averages. Annual averages do not capture hourly or marginal emissions, or the lifecycle emissions of generation.

the facility-level energy E_{facility} of equation 15.9, which already includes IT support and PUE once, scaled by the electricity grid's carbon intensity $\text{CO}_{2,\text{grid}}$:

$$\text{CO}_{2,\text{operational}} = E_{\text{facility}} \times \text{CO}_{2,\text{grid}} \tag{15.10}$$

Applied to a mid-scale training run, the equation separates the facility energy from the grid factor that converts it into emissions. The worked account below builds that facility energy step by step from accelerator draw, rack support, and PUE before comparing the resulting emissions across two regional grids.

Napkin Math 15.5: Training emissions calculation

Problem: Quantify the operational carbon footprint of training a 7B parameter model on 64 A100 GPUs for 14 days, assuming continuous accelerator draw at the rated power and the stated rack-support and facility overheads.

Math:

1. Compute energy: 64 GPUs $\times$ 400 W $\times$ 336 hours = 8,601.6 kWh.
2. Apply IT support and PUE: host processors, memory, and networking add to the accelerator load. This calculation borrows the ratio of IT power to accelerator power from the DGX H100 rack profile itemized later in table 15.7. That ratio is a modeling proxy for an A100 server's support load rather than a measurement of it. It scales accelerator energy by 1.38 to 11,827.2 kWh, so total facility energy = 11,827.2 kWh $\times$ 1.12 = 13,246.5 kWh.
3. Calculate emissions: On US average grid (429 g/kWh CO_2), emissions = 5.7 t. On Quebec hydro grid (20 g/kWh), emissions = 264.9 kg.

Systems insight: The grid factor alone produces a 21.4× difference in training emissions for this pairing, the same placement lever quantified in notebook 15.2. The model architecture is unchanged.

15.5.5 Embodied carbon assessment

Moving the 64-A100 training run of the preceding section from the US-average grid to Quebec's hydro-dominated grid cuts operational emissions by more than an order of magnitude, yet the carbon emitted to fabricate those 64 GPUs in the semiconductor foundry does not fall by a single gram. The training-emissions notebook priced electricity only, where a lower grid emissions factor reduces operational carbon and raises the amortized manufacturing share from about 1.6 percent of total emissions on the US-average grid to about 25.8 percent on Quebec's grid. Whether embodied carbon remains a minor correction or becomes a major share depends on the manufacturing footprint, utilization, and service life, so those quantities must enter the comparison before an engineer can choose between reducing electricity use and extending hardware service.

Embodied carbon encompasses emissions from raw material extraction, semiconductor fabrication, assembly, transportation, and end-of-life disposal. For AI hardware, the energy-intensive nature of advanced semiconductor processes dominates manufacturing emissions, including wafer fabrication at advanced process nodes, high-bandwidth memory production, and packaging. An NVIDIA A100 GPU carries an embodied footprint of approximately 150 kg CO_2e per unit (Luccioni et al. 2023). NVIDIA's HGX H100 product carbon footprint implies roughly 164 kg CO_2e per H100 when the baseboard footprint is allocated evenly across its eight GPUs (NVIDIA Corporation 2025). Equation 15.11 amortizes this embodied carbon over the hardware lifetime to compute per-use emissions:

$$\text{CO}_{2,\text{embodied,daily}} = \frac{\text{CO}_{2,\text{manufacturing}}}{T_{\text{lifetime}} \times 365} \tag{15.11}$$

with T_{lifetime} in years.

The manufacturing emissions have already occurred when the hardware enters service. Amortization allocates that burden over time; it does not describe additional manufacturing emissions accumulating each day. Utilization and service life determine how much useful work shares the initial burden.

Systems Perspective 15.1: Embodied carbon amortization

Renting a GPU for an hour uses hardware whose manufacturing emissions have already occurred. The workload account allocates a share of those emissions to the rental interval, regardless of whether the monetary fee explicitly prices carbon.

Expressing equation 15.11 per hour and multiplying by the job's duration makes that allocation explicit:

$$\text{CO}_{2,\text{total}} = \text{CO}_{2,\text{operational}} + \left(\frac{\text{CO}_{2,\text{manufacturing}}}{T_{\text{lifetime}} \times 365 \times 24} \times T_{\text{job}} \right)$$

with T_{lifetime} in years and T_{job} in hours.

Scenario: Training a model for 10 hours on 8 NVIDIA H100s. This accelerator-only illustration uses TDP as a constant-power assumption and excludes host, network, and facility energy. Its embodied share uses the allocated baseboard footprint; it is not a complete fleet lifecycle inventory.

- Operational: 8 × 0.7 kW × 10 hours = 56 kWh. At 0.429 kg/kWh (US-average grid) = 24 kg.
- Embodied: 8 × 164 kg = 1312 kg.
- Amortization: Lifetime = 3 years, the standard service life (26280 hours).
  - Hourly "Rent" = 1312 kg / 26280 hours ≈ 0.050 kg/h.
  - Job Cost = 0.050 kg/h × 10 hours = 0.5 kg.

Systems insight: For long-lived hardware in dirty grids, electricity dominates (24 kg vs. 0.5 kg). However, in clean grids (hydro, 0.020 kg/kWh), operational drops to 1.1 kg, making embodied carbon a significant fraction (~30.8 percent) of the total footprint.

The perspective's H100 job, like the lifecycle notebook later in this chapter, amortizes over the 3-year standard service life. The A100 calculation below keeps its own setup's 4-year extended life, which spreads the same manufacturing carbon over more service. For an accelerator with 150 kg embodied carbon (the A100 estimate cited earlier) and that 4-year data center lifetime, the first step is daily amortization: 150 kg / (4 years $\times$ 365 d/year) $\approx$ 0.103 kg/day. The second step assigns that daily share to the job: a training run lasting 14 days on 64 accelerators carries 64 $\times$ 14 days $\times$ 0.103 kg $\approx$ 92.1 kg CO_2 of amortized embodied carbon.

Adding that 92.1 kg embodied contribution to the 5,682.7 kg of operational emissions calculated earlier for the US-average grid yields the 1.6 percent manufacturing share noted at the start of this section, whereas pairing the same hardware burden with Quebec's 264.9 kg of operational emissions raises the manufacturing share to 25.8 percent. The manufacturing term did not change; the cleaner grid shrank everything around it.

15.5.6 Lifecycle carbon accounting

Reporting only the final training run's electricity and the serving cluster's power bill leaves out the failed architecture sweeps that preceded the run and the manufacturing emissions of the racks that host it. A complete workload account (equation 15.12) groups operational emissions into training and inference, adds embodied emissions, and requires that development experiments and supporting services enter the inventory rather than disappear when accounting starts at the final run:

$$CO_{2,\text{lifecycle}} = CO_{2,\text{training}} + CO_{2,\text{inference}} + CO_{2,\text{embodied}} \tag{15.12}$$

Each term has its own physical drivers (figure 15.8). The embodied term $CO_{2,\text{embodied}}$ is fixed before the hardware runs, at roughly 150 kg to 164 kg CO_2e per accelerator for the devices priced above. The training term $CO_{2,\text{training}}$ is a one-time cost set by cluster scale, PUE, and checkpoint I/O, while the inference term $CO_{2,\text{inference}}$ recurs for as long as request traffic and KV-cache residency keep serving replicas powered. Their relative proportions therefore shift over the deployment lifetime, as the worked account in notebook 15.6 and the serving crossover in section 15.7.3 show.

The equations become a decision tool once each factor is measured or declared. A worked account for a large training run shows how the operational and embodied terms combine, and which of them dominates.

Napkin Math 15.6: Partial lifecycle carbon estimation

Problem: Estimate facility operational emissions plus allocated accelerator-baseboard manufacturing emissions for training a 70B parameter model.

Variables: 2,048 H100 GPUs, 30 days, 700 W TDP, rack-profile support power, PUE 1.12, grid intensity 429 g CO_2/kWh.

Math: Accelerator power = 2,048 $\times$ 0.7 kW $\approx$ 1,433.6 kW. Applying the DGX H100 rack profile for host, memory, networking, and conversion support raises the IT load to 1,971.2 kW; facility power = 1,971.2 kW $\times$ 1.12 $\approx$ 2,207.7 kW. Energy = 2,207.7 kW $\times$ 24 h/day $\times$ 30 days $\approx$ 1,589,575.7 kWh. Emissions $\approx$ 681.9 t CO_2.

Embodied: Assume manufacturing footprint is $\approx$ 164 kg CO_2 per H100 GPU, allocating NVIDIA's HGX H100 baseboard product carbon footprint evenly across its eight GPUs (NVIDIA Corporation 2025). This worked account uses time-based allocation: it assigns the job the fraction of manufacturing emissions corresponding to its 1-month occupancy of a 3-year service life. Under that declared rule, (2,048 $\times$ 164 kg) $\times$ (1 month / 36 months) $\approx$ 9.3 t CO_2.

Result: 681.9 t + 9.3 t $\approx$ 691.3 t CO_2.

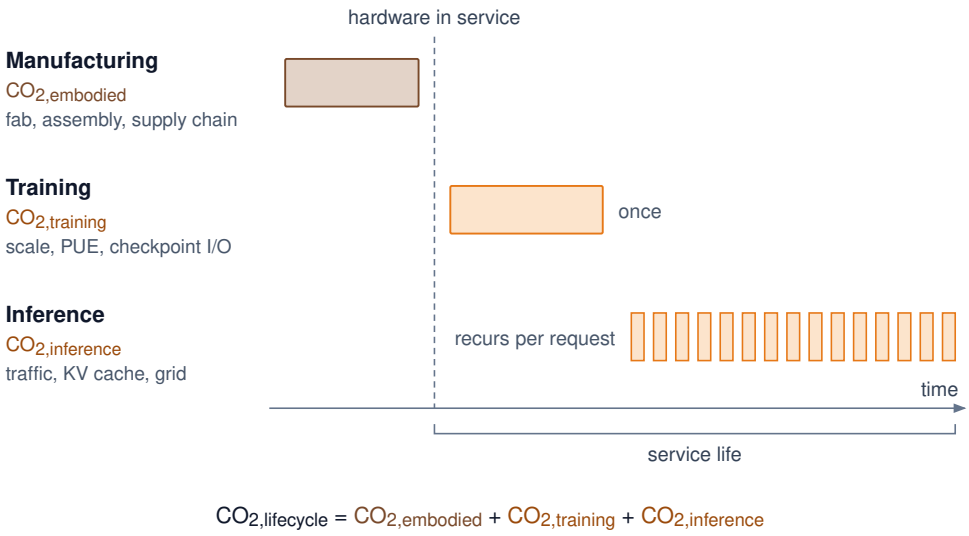

Figure 15.8: The Carbon Lifecycle of an ML System: Manufacturing contributes embodied carbon, while training and inference contribute operational carbon through energy, facility overhead, and grid intensity. Block widths show timing, not magnitude. The relative size of each term is workload-specific: serving volume, grid factor, and hardware lifetime determine which one dominates.

Systems insight: The operational term dominates this partial lifecycle account. It includes facility electricity and allocated HGX baseboard manufacturing, but omits manufacturing of host systems, networking, storage, and facility infrastructure, as well as other unmodeled lifecycle stages. A complete footprint must extend the inventory to those boundaries.

Time is one defensible allocation basis, not a universal one. Shared fleets may instead allocate embodied emissions by accelerator-hours, reserved capacity, or a workload's share of useful computation. Google's cloud carbon reporting combines two of these bases, allocating the energy and emissions of shared machines by each customer's resource reservations and hourly measured usage (Schneider and Mattia 2024). The chosen denominator must be reported and applied consistently; otherwise a low-utilization job can appear artificially inexpensive merely because another tenant absorbs the idle capacity. These workload shares support engineering comparisons; they do not determine the reporting year or category in a corporate emissions inventory. Procurement decisions need the unallocated fleet total as well as the workload-attributed share, because amortization changes ownership of an emission in the inventory but does not remove it from the atmosphere.

Every lifecycle estimate should therefore carry an uncertainty interval rather than a single authoritative total. Average workload power, PUE, grid intensity, hardware lifetime, utilization, and manufacturing footprint are measured or modeled with different errors. A sensitivity analysis varies the assumptions that could reverse the decision, including clean-grid versus carbon-intensive operation, three- versus five-year service life, and low versus high serving volume. It then reports whether the preferred design remains preferred. Precision in arithmetic cannot compensate for an unstable system boundary.

The balance among these terms changes with serving volume and duration. A high-volume service can exceed its entire training footprint within months, or within days at higher traffic, a crossover that section 15.7.3 works through for a deployed model.

The account becomes an engineering tool when it runs on every job rather than once after the fact. A partial lifecycle carbon calculator (listing 15.1) combines measured energy, the rack support multiplier, a declared facility PUE, an explicitly selected grid-intensity factor, and amortized embodied

carbon. Its example reproduces the 64-A100 run, with each input drawn from the facility, grid, or hardware profile it describes; section 15.9.4.1 compares the sources of the grid factor.

Listing 15.1: Lifecycle Carbon Calculation: Computing total carbon footprint including operational and embodied emissions.

```python
from mlsysim import Hardware, Infrastructure, ReferenceStats, Systems


def calculate_carbon_footprint(
    gpu_power_watts: float,
    num_gpus: int,
    training_hours: float,
    it_power_multiplier: float,
    pue: float,
    grid_intensity_gco2_kwh: float,
    gpu_embodied_kg: float,
    gpu_lifetime_years: float,
) -> dict:
    """Estimate operational emissions plus allocated GPU manufacturing emissions."""

    # Operational emissions (plain floats: watts, hours, kWh, grams)
    gpu_energy_kwh = gpu_power_watts * num_gpus * training_hours / 1000
    it_energy_kwh = gpu_energy_kwh * it_power_multiplier  # + host, net
    facility_energy_kwh = it_energy_kwh * pue  # overhead, applied once
    operational_kg = (
        facility_energy_kwh * grid_intensity_gco2_kwh / 1000
    )

    # Embodied emissions (amortized by time over the service life)
    daily_embodied = gpu_embodied_kg / (gpu_lifetime_years * 365)
    training_days = training_hours / 24
    embodied_kg = num_gpus * training_days * daily_embodied

    return {
        "energy_kwh": facility_energy_kwh,
        "operational_carbon_kg": operational_kg,
        "embodied_carbon_kg": embodied_kg,
        "total_carbon_kg": operational_kg + embodied_kg,
        "embodied_fraction": embodied_kg
        / (operational_kg + embodied_kg),
    }


# Example: 7B model training on a 64-GPU A100 cluster for 14 days.
# Each input names its source: PUE is a facility property, the
# grid factor a property of the electricity supply.
_rack = Systems.Racks.DGX_H100_4Node  # proxy support-load profile
_facility = Infrastructure.FacilityCooling.BestAir  # PUE 1.12
_grid = Infrastructure.Grids.US_Avg  # annual-average factor
_run = ReferenceStats.MidScaleTrainingCarbonProfile
result = calculate_carbon_footprint(
    gpu_power_watts=Hardware.Cloud.A100.tdp.m_as("W"),  # 400 W
    num_gpus=Systems.Clusters.Lab_64_A100.total_accelerators,  # 64
    training_hours=_run.TrainingDays.m_as("h"),  # 336 h
    it_power_multiplier=(
        _rack.it_power / _rack.accelerator_power
    ).m_as(""),
    pue=_facility.pue,
    grid_intensity_gco2_kwh=_grid.carbon_intensity_g_kwh,
    gpu_embodied_kg=Hardware.Cloud.A100.embodied_carbon_kg,
    gpu_lifetime_years=(  # 4-year extended service life
        Infrastructure.Pricing.Capital
        .AcceleratorExtendedServiceLifeYears.rate
    ),
)
print(
    f"Total carbon footprint: {result['total_carbon_kg']:.0f} kg CO2"
)
print(f"Embodied fraction: {result['embodied_fraction']:.1%}")
```

Teams can integrate this partial carbon account directly into orchestration dashboards. Calculating operational and embodied emissions for individual training runs, however, captures only one dimension of fleet resource use. At campus scale, data centers introduce power-delivery and cooling constraints that individual run metrics leave invisible.

15.6 Data Center Energy and Resource Consumption

A training cluster can fit its accelerator power budget yet exceed its facility budget once support systems and cooling enter the account. Whether a facility can host the cluster at all is a feasibility question; the sustainability question is whether a change in power delivery or cooling reduces lifecycle cost while preserving that feasible operating envelope. The relevant comparison includes annual energy, local water use, equipment lifetime, and the energy wasted when cooling failures interrupt useful work.

Facility sustainability therefore has to be read as a chain of constraints rather than as a single PUE number. Persistent megawatt demand sets the grid and emissions exposure. Power delivery determines how much electricity becomes useful IT load. Cooling sets the facility overhead and the water drawn to remove heat. Water use then determines whether a technically efficient site is locally sustainable. The physical mechanisms below matter here because each assigns losses or resource use to a different accounting boundary.

15.6.1 Data center energy and AI workloads

Even when an engineer halves the energy of a single model, campus electricity demand still climbs if new data-center capacity and serving traffic expand faster than per-job efficiency. At facility scale, the optimization target is no longer a single model in isolation but the aggregate demand, facility overhead, and grid context surrounding every watt of IT power. The same IEA *Energy and AI* analysis behind the 2030 electricity projection cited at the chapter's opening also tracks emissions, a different quantity. It estimated emissions from data-center electricity use at about 180 Mt at the time of the analysis and projected roughly 300 Mt by 2035 in its Base Case, still below 1.5 percent of total energy-sector emissions (International Energy Agency 2025).[12] The energy burden grows with data-center capacity, training workloads, and inference demand (Patterson et al. 2021), so component efficiency must be evaluated against fleet expansion (Thompson et al. 2021; Dodge et al. 2022).

The relevant facility quantity is persistent megawatt-class load, not the model name alone. Hyperscale campuses house large fleets of accelerator servers and hold loads of 100 MW or more for years at a time.[13] A newer accelerator lowers a job's energy only when that workload realizes the efficiency gain. Even then, lower energy per job need not reduce campus demand if the fleet uses the released capacity to perform more work.

Whether that demand keeps growing depends on the race between workload growth and efficiency. Masanet et al. (2020) show that efficiency gains held global data-center energy nearly flat through the 2010s even as computing output grew several-fold. Sustained AI demand growth can erode that offset, which is why the IEA projection cited earlier rises rather than remaining flat.

Beyond computational demands, cooling can account for 30–40 percent of facility energy in the conventional air-cooled data-center designs surveyed by Ebrahimi et al. (2014). This is a design-regime range, not a constant for modern hyperscale or liquid-cooled facilities; section 15.6.3 derives the site-specific overhead through PUE and the cooling architecture.

The United States alone illustrates how cloud and AI infrastructure can reshape national energy planning. US data center electricity consumption more than tripled from 58 TWh in 2014 to 192 TWh in 2024, reaching 4.7 percent of national supply, according to the Lawrence Berkeley National Laboratory (LBNL) estimates that figure 15.9 plots (Shehabi et al. 2024; Smith et al. 2026). LBNL's reference case projects 649 TWh for 2030, about 11.8 percent of US electricity, inside a scenario range of 521 to 843 TWh (Smith et al. 2026). These scenarios describe possible demand, not a measured capacity ceiling. They motivate evaluating efficiency gains alongside the generation, transmission, and connection capacity needed to serve that demand.

Every kilowatt-hour in those national totals begins as electricity delivered to a rack and leaves as heat that must be removed from it. Delivering 120 kW to a single cabinet (the rated draw of one

[12] **Data center emissions scale**: Those emissions come from electricity demand that IEA's 2025 *Energy and AI* analysis tracks separately. Data centers consumed about 1.5 percent of global electricity (415 TWh) in 2024, a share projected to approach 3 percent by 2030, with AI contributing to demand growth (International Energy Agency 2025). The largest hyperscale facilities can draw over 100 MW continuously, equivalent to powering tens of thousands of homes.

[13] **Hyperscale data center footprint**: A hyperscale campus makes long-lived investments in substations and transmission connections. Algorithmic improvements can reduce demand but do not remove the costs already incurred for that capacity. Supply contracts and operating schedules remain adjustable, subject to their terms and the physical connection.

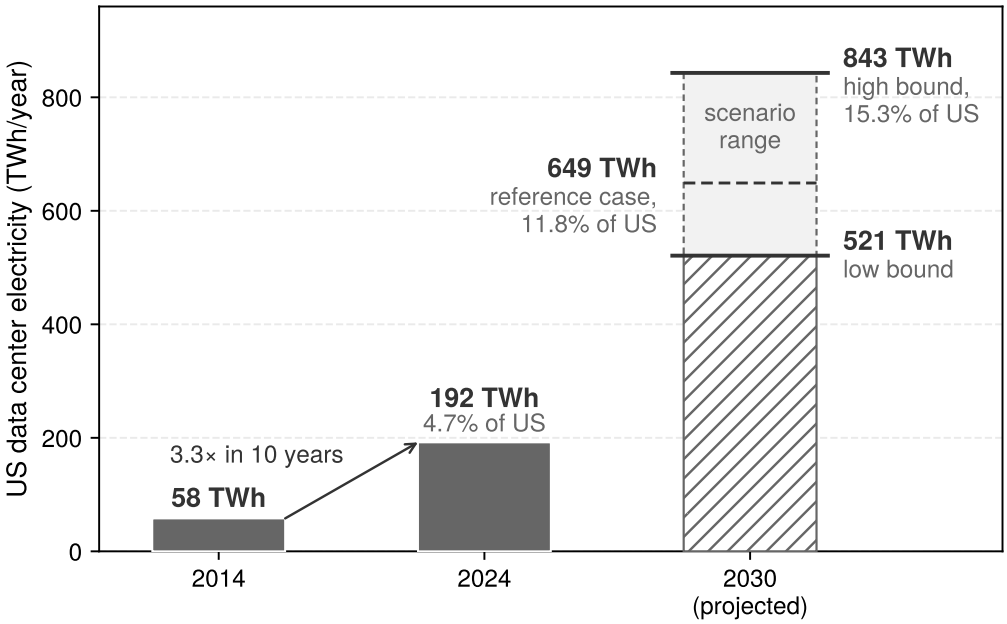

Figure 15.9: The Energy Gap: US data center electricity consumption more than tripled from 58 TWh (2014) to 192 TWh (2024). LBNL projects 649 TWh by 2030 in its reference case, within a range of 521 to 843 TWh (Shehabi et al. 2024; Smith et al. 2026), with AI workloads treated as an important driver of growth. At the high end, data centers would consume approximately 15.3 percent of US electricity, motivating evaluation of the generation, transmission, and connection capacity needed to serve that projected demand.

GB200 NVL72 rack) starts at the substation, and every conversion stage between the grid and the chip takes its share along the way.

15.6.2 Power delivery

Capacity planning fails when engineers budget only accelerator TDP, because transformers, uninterruptible power supplies, rack distribution, and voltage regulators each dissipate part of the incoming grid power as heat before a single watt reaches the silicon. Section 2.6.1 traced that electrical path from the utility feed to the voltage regulator, and equation 15.2 measures the facility overhead along it. For a sustainability account, the question is where along that path energy and reliability margin are lost, and which boundary carries each loss. The first constraint sits outside the building: large grid-supplied ML facilities need sufficient substation and utility-interconnection capacity. Internal distribution improvements cannot raise that import limit, and on-site generation or storage adds a separate supply whose capacity, duration, and reliability must enter the design.

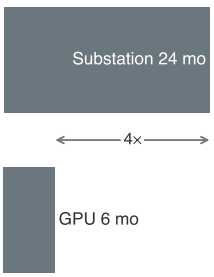

Substation lead time can run 4× longer than GPU procurement.

Systems Perspective 15.2: The interconnection queue

Deploying 10,000 GPUs and deploying the electrical substation required to power their IT load (9.6 MW, after the 7 MW GPU subtotal is expanded for rack support power) operate on different timelines. In this planning comparison, GPU procurement takes 6 months, while permitting, engineering, procurement, construction, and grid connection for a new 10+ MW substation take 24 months. The infrastructure path is therefore 4× longer than the accelerator path. Actual schedules depend on the site, utility, equipment availability, and approvals, so the ratio illustrates a dependency rather than a universal lead time.

Systems insight: In this comparison, power delivery determines when the fleet can begin useful work even if the accelerators arrive on schedule. The US generation-and-storage interconnection queue is distinct from a data center's load-connection request; neither queue total directly

measures available capacity for a new ML facility. Procurement must therefore follow a site-specific power-delivery schedule. Otherwise, delivered silicon can remain "electrically stranded" while the facility waits for its connection or transformer.

Inside the facility, every conversion stage turns part of the delivered power into heat. Along the path traced in section 2.6.1 (figure 2.12), the utility transformer loses about 1–2 percent, an online double-conversion UPS 3–5 percent (or 1–2 percent in eco mode), the power distribution unit (PDU) 2–3 percent, and baseboard voltage regulation 5–8 percent. Compounded, these stages dissipate roughly 9–17 percent of grid power before it reaches a transistor, and server power supplies add their own conversion loss on top. The accounting boundary decides where each loss appears. Losses upstream of the IT boundary, such as the transformer, UPS, and facility distribution, raise PUE. Server power supplies and baseboard voltage regulators sit inside the measured IT load, and their heat must be removed by the same cooling system that serves the accelerators.

Higher-voltage distribution is one of the few design changes that shrinks this overhead directly, because the same power at a higher voltage needs less current, and resistive distribution losses I^2R fall with the square of that current. The 48 V rack distribution of section 2.6.1.1 applies that principle, and its realized saving depends on conductor design and the added conversion stages.

At ML rack power densities, the difference between 95 percent and 98 percent distribution efficiency is meaningful. Using 33 kW as a reference rack load, a first-order estimate applies the difference between those efficiencies as 3 percent of that load, yielding approximately 0.99 kW per rack. Across 300 racks, the estimated reduction is 297 kW, equivalent to roughly 36 GPU nodes or 9 four-node racks under this power mix. Sustaining that reduction over a 3-year lifecycle at \$0.07/kWh gives estimated electricity savings of ~\$546,361. An exact fixed-load comparison divides delivered power by each efficiency before subtracting input powers. The released capacity can support more work or lower consumption; spending it on additional nodes does not preserve the same electricity saving.

Section 2.6 sizes the power envelope of a four-node DGX H100 rack; the same rack also shows which parts of a budget are IT load and which are facility overhead. A rack-attributed power budget itemizes more than GPU TDP (table 15.7). In this scenario, host, fabric, and equipment-side conversion losses belong to the IT load and sum to the IT subtotal; the row below it is an illustrative non-IT allocation at the stated PUE. A facility calculation multiplies the IT subtotal by its own PUE rather than reusing that allocation, and facility-side conversion losses are covered by PUE and must not be added again.

Table 15.7: Power Budget for a Four-Node DGX H100 Rack: GPUs consume two thirds of this illustrative rack-attributed total. Host systems, networking, and equipment-side conversion losses form the remaining IT load; the final overhead row represents the scenario's non-IT allocation. The conversion row excludes facility-side losses, which belong in PUE. Applying a different facility PUE requires replacing the non-IT allocation, not multiplying the full table total.

Component	Power (kW)	% of rack total
GPU compute	22.4 kW	67%
Host CPUs and DRAM	3.2 kW	10%
NVSwitch fabric	1.6 kW	5%
InfiniBand HCAs	0.8 kW	2%
IT-equipment conversion losses	2.8 kW	8%
IT subtotal	**30.8 kW**	
Allocated non-IT overhead (PUE ~1.09)	2.7 kW	8%
Total	**33.5 kW**	**100%**

Evaluating DGX H100 power allocation (table 15.7) reveals that 32 GPUs at 700 W each deliver 22.4 kW. Host, fabric, and conversion losses, however, raise the IT subtotal to 30.8 kW, about 38 percent above the GPU subtotal. That IT subtotal is the load to multiply by a facility's PUE; the table's final row is one such allocation, which brings the rack-attributed total to about 50 percent above the GPU subtotal. Infrastructure planners who size their facility based on GPU TDP alone will underestimate the electrical load and may discover during commissioning that their power capacity is insufficient.

Synchronized training also changes the shape of the load. The correlated power swings introduced with grid dynamics (section 15.3.2) now appear inside the delivery chain, where coordinated computation and communication phases align across thousands of accelerators. Section 2.6.1.1 analyzes how decoupling capacitors, voltage regulators, and UPS controls filter these **synchronous transients** before they reach the utility connection. The same section estimates the aggregate ramp, in gigawatts per second, when a training cluster steps from communication to computation in unison. That estimate describes a load-side scenario rather than a measured utility input, and its sustainability consequence appears at the interconnection. Rack- or facility-level energy buffering (section 15.3.2), sized to the measured energy deficit, the energy a step draws above what the utility connection supplies during the ramp, rather than to a fixed allowance per rack, lets the grid see a smoother load. A well-characterized, controllable load is easier both to connect and to schedule against grid conditions.

Checkpoint 15.3: Power delivery physics

Verify the data center power path:

- ☐ **Load correlation**: explain why synchronous ML training produces more correlated load changes at the facility connection than independent web traffic.
- ☐ **Loss accounting**: identify which conversion losses in the delivery chain raise PUE and which remain inside the measured IT load.
- ☐ **Power budgeting**: use table 15.7 to explain why a node with 8 GPUs at 700 W each must be budgeted well above its 5.6 kW GPU subtotal.
- ☐ **AC-DC conversion**: identify the conversion-stage efficiencies, distribution losses, load profile, and reliability requirements needed to compare rack-level and server-level conversion at the same delivered IT load.

At the largest scales, the data center's power draw represents a significant fraction of the local electrical grid's capacity. A 100,000-GPU cluster, roughly a hundred times the book's 1,024-GPU reference cluster, at 700 W per GPU has a 70 MW accelerator subtotal. Rack-profile support power, however, raises the IT load to 96.2 MW before PUE. With PUE, the total facility draw approaches 107.8 MW, the load of a small city. Such installations require dedicated feeds from the electrical grid, often with purpose-built substations and transmission lines. The substation lead time of 24 months in the interconnection-queue perspective above makes power availability one of the longest-lead-time constraints in building ML infrastructure.

At large-cluster scale, the same delivery limit becomes a grid-procurement problem rather than an electricity bill. The problem worsens because compute-optimal allocation (section 1.3.2) grows parameters and training tokens together, so training work, which scales with their product, grows roughly with the square of model size, and energy follows it unless realized FLOPs per joule improve.

The electricity consumed by a single large-model training campaign can become nonnegligible relative to the output of a power plant. A 100 MW training facility operating at full capacity for one year consumes 876 GWh. A 100 MW wind farm cannot match that figure (at a 30 percent capacity factor, the share of its rated output it delivers on average, a wind farm produces about 262.8 GWh per year, so the training facility would require the equivalent of approximately 3.3 large wind farms). Once power is available, its carbon intensity determines the emissions attributed to each megawatt-hour under the chosen accounting method, and at facility scale that factor becomes a siting decision that persists for the life of the building. Low-carbon electricity can coincide with favorable prices, but price and emissions factor are separate tests that a siting comparison must evaluate alongside power availability and the workload's location constraints.

15.6.3 Cooling energy, heat rejection, and water

Nearly every watt of electrical power that enters a server cabinet leaves as heat, so a dense rack acts as a continuous heater whose thermal output must be removed at the rate it consumes electricity or the silicon throttles and shuts down. Section 2.6.2 established how air and liquid loops carry that heat and which cooling class a given rack density requires; for a sustainability account, the question

is what rejecting that heat to the outside environment costs in additional facility electricity and local water.

PUE (equation 15.2) captures facility fans, chillers, and facility-side conversion losses as overhead on top of the IT load. If facility cooling consumes 3–8 percent of IT power, that contribution alone adds 0.03–0.08 to PUE, and facility-side electrical losses and other non-IT loads add to it. Server fans, by contrast, lie within the IT load itself, so their energy sits in the denominator as well as the facility total, and reducing server-fan power changes the IT load rather than producing the same PUE reduction. Cooling choices must therefore be compared using total facility energy for the same useful work, not PUE alone.

The cooling class of section 2.6.2 settles thermodynamic feasibility, that the proposed cooling medium maintains die temperatures and component serviceability under peak workload. Lifecycle economics then compares installed capital expenditure, pumping and fan parasitics, water consumption, maintenance overhead, and equipment lifespan under the same computational load. A higher-density liquid-cooled installation may reduce operational energy, but its capital payback period cannot be inferred from rack power alone.

Napkin Math 15.7: The cooling overhead

Consider a 1,024-GPU cluster of H100s at 700 W each. To isolate the cooling choice, this simplified calculation treats accelerator power as the entire IT load and assigns all non-IT overhead to cooling; a procurement estimate must restore host, memory, network, and storage support before applying PUE.

- IT power: $1{,}024 \times 700$ W = 716.8 kW
- Overhead removed: with the $(\text{PUE} - 1)$ overhead arithmetic of notebook 15.3, moving from air cooling (PUE 1.55) to liquid cooling (PUE 1.08) removes 716.8 kW $\times$ 0.47 = 336.9 kW of continuous power.

Savings: At \$0.07/kWh, the annual savings are approximately \$206,585. Over a 3-year hardware lifecycle, the cooling savings alone total \$619,754, which is the figure to set against the installed capital cost of the liquid-cooling system when computing its payback period.

The final link in the cooling chain is **heat rejection**: getting heat from the facility water loop to the outside environment. A cooling tower sprays warm water over a fill medium, packing that spreads the water into thin films to enlarge its contact with the air, and uses evaporation to carry heat into the atmosphere. Evaporating a kilogram of water removes its latent heat of vaporization, 2,260 kJ/kg. Warming the water without evaporating it, called sensible heating, instead depends on both the specific heat capacity, 4.18 kJ/kg/K, and the permitted temperature rise. Evaporation can therefore carry substantial heat with a small mass flow, but the evaporated water must be replaced.

A 10 MW data center with evaporative cooling towers can consume well over 100 million liters of water per year, depending on climate and cooling design, which has become a significant concern in water-stressed regions. At the evaporative WUE of 1.8 L/kWh used earlier, continuous operation consumes about 158 million liters per year, roughly 350–415 household-years of water use under the earlier annual household baseline. As ML data centers grow to 100 MW and beyond, their water footprint becomes a meaningful factor in local resource planning.

Dry coolers use fans to blow air over a radiator without evaporating cooling water. They avoid that direct water consumption, but the added fan electricity carries its own water burden at the generating plant, in a different watershed from the facility (figure 15.10). Heat rejection depends on the temperature difference between coolant and ambient air, so hot weather can reduce capacity or require more fan power. Hybrid designs use dry cooling when conditions permit and evaporative assistance when needed, trading water consumption against electrical demand and thermal margin.

Waste heat reuse treats thermal output as a resource. A **district heating** connection supplies a municipal heat network, while **industrial coupling** directs recoverable heat toward a compatible process. The benefit depends on the temperature, timing, and distance of the receiving demand, along with any pumping or heat-pump energy needed to supply it. Useful recovery can displace another heat source, although it does not eliminate the fleet's electrical consumption[14] (Ebrahimi

[14] **PUE boundary**: The facility-to-IT ratio glossed with the opening energy boundary (The Green Grid 2007), formalized in section 15.4.1.3, says nothing about where recovered heat goes. Lowering non-IT overhead improves facility efficiency, but the metric does not capture the carbon intensity of electricity or the water consumed for cooling. Carbon usage effectiveness and water usage effectiveness therefore expose constraints that PUE alone leaves invisible.

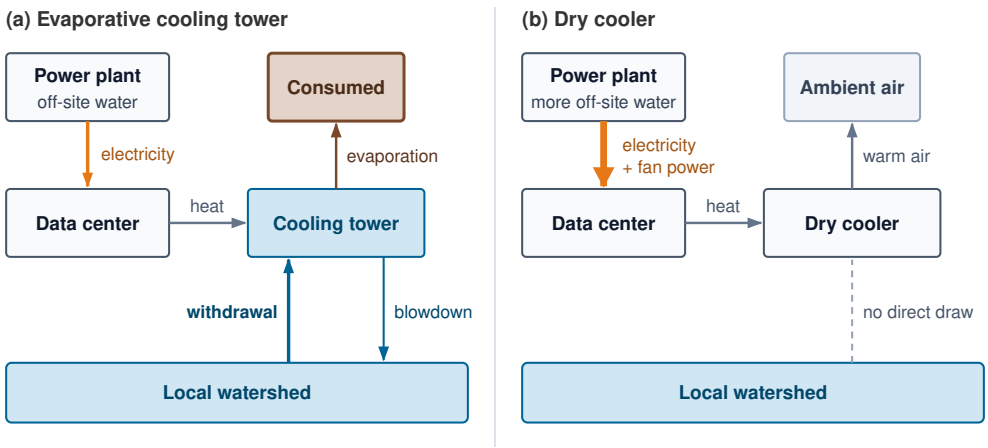

Figure 15.10: Direct and Indirect Water Accounts: An evaporative cooling tower (a) rejects heat by consuming water withdrawn from the local watershed, returning only its blowdown. A dry cooler (b) draws no local cooling water but spends more electricity on fans, and generating that electricity consumes water in the power plant's watershed. Water is therefore a separate, location-weighted account traded against energy. The schematic shows which flows exist, not their sizes.

et al. 2014). A liquid-cooling loop returning water at 50–65 degrees Celsius supplies low-grade heat that may serve nearby district heating, greenhouses, or other suitable sinks. Usability depends on the required supply temperature, seasonal demand, and transport losses. Because municipal heating networks often require supply temperatures above 70 degrees Celsius, the facility may need an electric heat pump to raise the coolant temperature before delivery. If an illustrative 10 MW thermal stream has an 80–90 percent recovery fraction, 8–9 MW is recovered before accounting for heat-pump electricity. Heat-pump output includes that electrical input and must be accounted for separately.

Crediting useful heat does not reduce PUE below 1.0 because PUE compares total facility electrical input against IT electrical input, and exporting heat downstream does not reduce the electricity the building draws from the grid. Exported heat is instead reported through energy reuse effectiveness (ERE), which subtracts the reused energy from total facility energy before dividing by IT energy, or through a separate net-energy balance.

The viability of waste heat reuse depends on the proximity of heat consumers. Urban data centers, despite their higher land and electricity costs, are often better positioned for waste heat reuse than remote facilities because they are close to residential and commercial heating loads. The result is a counterintuitive economic optimization: a data center in a Nordic city may have higher electricity costs but lower net operating costs after waste heat revenues, compared to a remote facility with cheaper electricity but no heat consumers nearby.

A cooling failure also costs energy, because the run repeats the work done since its last checkpoint. Rerunning 30 to 60 minutes on all 10,000 accelerators of a training cluster at rated power repeats about 3.5–7 MWh of accelerator energy, before host and facility overhead. Power drawn by idle accelerators during the outage adds to that total. ML goodput, which counts only useful fleet output, discounts the repeated work, but the energy account still pays for it. Section 2.6.4 develops the cooling failure modes and their blast radius, and section 7.8 develops the checkpoint recovery that bounds the rollback. Together, power-delivery losses and heat rejection set the facility overhead on every kilowatt-hour entering the building; how that electricity budget divides between one-time model training and continuous production serving determines where workload-level optimization must focus next.

↩ ML goodput, the share of fleet output that counts as useful work, is established in section 9.2.

15.7 Training vs. Inference Energy Analysis

Training a large language model is a concentrated, highly visible energy event that consumes thousands of megawatt-hours over weeks or months. Serving the same model to a billion daily queries is the opposite shape of load: energy spent continuously, query by query, year after year.

Training and inference share opportunities to reduce computation, data movement, and idle power, but their success criteria differ. Training must reach the target model quality within its completion budget; inference must serve the required demand at the agreed quality and latency. The lifecycle energy account identifies which phase contributes most, while those constraints determine which interventions can reduce its cost without sacrificing useful work.

15.7.1 Carbon benchmarks across AI workloads

Two published carbon estimates for the same neural architecture search, an automated search that trains many candidate architectures to find a good one, can differ by orders of magnitude when one assumes older GPUs, an average facility PUE, and a fossil-heavy grid (Strubell et al. 2019) while the other measures specialized accelerators in an efficient, low-carbon facility (Patterson et al. 2021).[15] Even though the IEA estimate cited earlier puts emissions from data-center electricity use at about 180 Mt, below 1.5 percent of total energy-sector emissions, workload-level carbon benchmarks govern where engineering teams invest optimization effort, and those benchmarks are only as reliable as the lifecycle boundaries behind them.

The GPT-3 calculation in notebook 15.1 put one training run at hundreds of trans-Atlantic passenger round trips. Such run-level comparisons convey scale, but they do not by themselves describe the footprint of a deployed service. The operational lesson is that carbon accounting must distinguish training from recurring inference over a declared observation period. Meta's recommendation workloads (RM-1 through RM-5) split operational emissions roughly evenly between training and inference, while inference dominates its language-model example (figure 15.11). The public baselines, including GPT-3 and BERT-NAS, report training only; their bars do not establish a lifetime training-to-inference ratio. For high-traffic Archetype B deployments (section 1.6.1.2), both recurring serving and model updates belong in the account before deciding which phase offers the larger reduction.

15.7.2 Training energy demands

Even the training-only bars in the comparison above represent massive, concentrated power events: as early as 2020, the supercomputer Microsoft built for OpenAI combined 285,000 CPU cores and 10,000 GPUs (Langston 2020) operating continuously for months, a dated reference point that subsequent frontier clusters have multiplied severalfold. Those intensive computational loads generate heat that cooling infrastructure must continuously remove, the facility overhead quantified in section 15.6.3, and reducing the draw requires co-optimizing hardware architecture, parallelism strategy, and algorithmic efficiency.

Every single-run training account in this chapter, from the GPT-3 run to the 70B model of notebook 15.6, also prices only the winning configuration. Development multiplies each run, by 5× for the sweep in section 15.5.2, while periodic retraining during service repeats part of that cost. Training energy is therefore paid once per completed run but recurs across the model's development and update cycle.

15.7.2.1 Distributed systems energy optimization

Distributing a training job across accelerators introduces coordination costs alongside the energy spent on computation. Distributed parallelism strategies exchange gradients, parameters, or activations while synchronization stalls leave idle accelerators drawing baseline power.[16] The resulting energy share cannot be inferred from cluster size alone because topology, collective algorithm, message volume, overlap, and utilization determine the account. Some AllToAll patterns also increase traffic faster than aggregate useful work as the participating set grows.

The Fleet Law (equation 1.3) turns this overhead into an energy account. If the fleet draws average power P_{avg} while executing one step on N accelerators, the energy of that step is approximately

$$E_{\text{step}} \approx P_{\text{avg}} \times T_{\text{step}}(N) = P_{\text{avg}} \left(\frac{T_{\text{compute}}}{N} + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}} \right) \tag{15.13}$$

The compute term accounts for local mathematical work; communication and synchronization make the distributed update possible. When these costs remain exposed, they extend the step while accelerators, hosts, and the fabric continue drawing power, adding powered time on the communication and coordination axes of C³. Their energy share depends on power during those

15 **Neural architecture search (NAS) carbon cost**: The ~284,000 kg CO_2e figure from Strubell et al. (2019) represents an unoptimized estimate for the search that produced the Evolved Transformer architecture, and it drops sharply when accounting for TPU hardware, grid mix, and the smaller proxy tasks on which the search evaluated candidates (Patterson et al. 2021). The systems lesson is boundary discipline: search, tuning, and failed trials belong inside the carbon accounting boundary, but reusable architecture search should be amortized across all downstream deployments.

16 **Parallelism energy overhead**: Data, model, and pipeline parallelism impose distinct communication and idle-power costs. Data parallelism reduces gradients, tensor parallelism exchanges partial activations, and pipeline parallelism introduces bubbles. A valid comparison reaches the same model-quality target under the same completion constraint, then measures each strategy's runtime and system energy, including communication and waiting. Fixing time-to-train by assumption would hide one of the effects the comparison needs to expose.

↩ Parallelism strategy selection and the communication cost of each strategy are established in section 6.2.1.

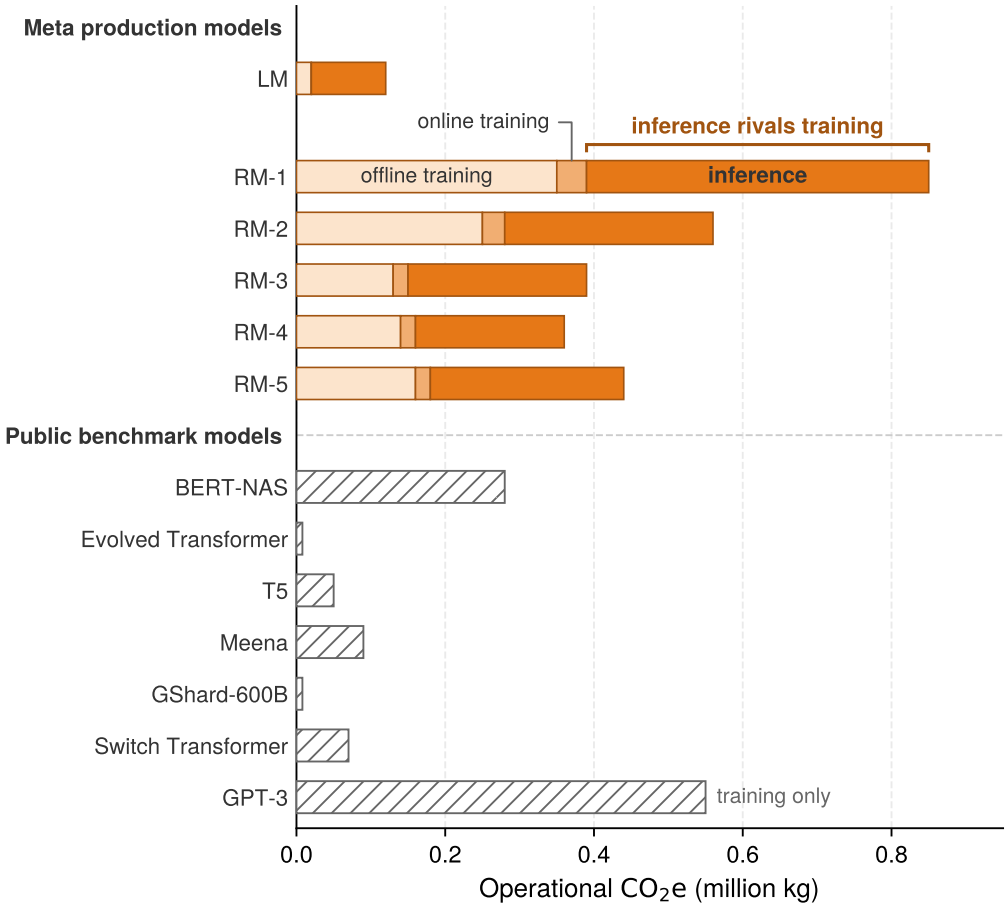

Figure 15.11: Operational Carbon Footprint of Large-Scale ML Tasks: Stacked bars show operational CO_2e over the reported accounting periods for Meta's recommendation models (RM-1 through RM-5), a language model (LM), and several open-source large models (BERT-NAS, Evolved Transformer, T5, Meena, GShard-600B, Switch Transformer, GPT-3). The recommendation-model bars divide roughly evenly between training (light and mid orange) and inference (solid orange), while the public open-source baselines report training-only footprints (hatched gray outline), showing that inference serving at scale can rival or exceed training emissions for deployed systems. Source: (Wu et al. 2022).

intervals as well as on their duration. A quarter of the step spent in unhidden communication accounts for a quarter of its energy only if average fleet power during that interval equals the average over the whole step. Accelerators stalled at a collective can draw less than at full arithmetic load, so the share must be read from a power trace. Overlap, smaller messages, and fewer stragglers shorten the step, and they save energy to the extent that the shorter powered interval is not offset by higher average power or by extra work to reach the same model quality. The parallelism choices of chapter 6 are therefore sustainability choices as well.

↩ The data-movement boundaries that set what each transferred byte costs, from on-package links through intra-node crossbars to inter-node switches, are established in section 5.2.

Beyond the communication term, coordinated resource management offers four levers with distinct effects on energy and emissions. At the cluster level, dynamic workload placement consolidates compatible jobs so unused hardware can enter lower-power sleep states. The savings must be verified at the facility meter rather than inferred from idle node counts. Carbon-aware scheduling (section 15.9.4) shifts flexible runs toward cleaner grid windows even when total energy use is unchanged, subject to data-movement overhead, completion deadlines, and reliability bounds. Inside active nodes, multi-tenant sharing raises accelerator compute and memory occupancy across concurrent tasks, though contention can lengthen execution and increase energy per completed job. Adaptive inference batching groups compatible requests to improve throughput within their latency

limits. Changing training batch size instead requires checking the total work needed to reach the target model quality. Each lever must therefore be evaluated against completed useful work and its service constraints rather than occupancy alone, because energy spent on work that ML goodput discounts, the scheduling, runtime, and program losses of section 9.2, still enters the energy account.

Energy-aware training frameworks reduce step energy directly rather than shifting where it is spent. Zeus improved training energy efficiency by 15.3–75.8 percent across its evaluated workloads by automatically finding energy-performance trade-offs (You et al. 2023), while Perseus reduced large-model training energy consumption by up to 30 percent by mitigating energy bloat, the energy spent running accelerators at full speed on work that must then wait for a slower pipeline stage or straggler (Chung et al. 2023). These tools make energy optimization available beyond hyperscale companies, yet all of these training costs stop accruing when the run stops. For high-volume services, the larger term is often the deployment that follows, as the crossover calculation in section 15.7.3 shows.

15.7.3 Inference energy costs

A model trained once can draw far more electricity over a year of high-volume serving than it ever consumed during training (Wu et al. 2022). Although a single inference request costs only a fraction of a watt-hour, serving executes every time the model answers a query, classifies an image, or scores a recommendation, scaling continuously with user traffic across search engines, recommendation systems, and generative AI services.

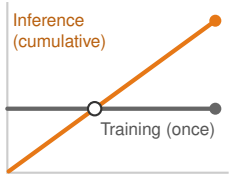

Cumulative inference emissions overtake the one-time training cost.

Consider a model serving 10 million queries per day at 0.001 kWh of attributed facility electricity per query, including its share of idle capacity and support systems. At this serving rate, daily energy consumption reaches 10,000 kWh (10 million queries $\times$ 0.001 kWh), which accumulates to 3,650,000 kWh over a year of continuous service. Evaluated against the average US grid carbon intensity of 0.429 kg/kWh, this electricity demand produces 1,565.9 t (1,565,850 kg) of greenhouse gas emissions annually.

At this scale, cumulative inference emissions exceed the 5.7 t of the earlier 7B/64-A100 training example, the smallest of the chapter's three training accounts and therefore the fastest to overtake, after approximately 1.3 days of deployment. The serving scenario and the training run describe different models, so the comparison does not trace one model's lifecycle. It shows how quickly this serving volume would overtake a training run of that size. For larger training runs, the crossover can shift to weeks or months depending on query volume, per-query energy, and grid intensity. The lifecycle perspective therefore sets the priority: optimize inference efficiency for widely deployed models, and focus training efficiency efforts on models that undergo frequent retraining or experimental iteration.

Production measurements in figure 15.11 show why this crossover must remain workload-specific. Inference dominates the language-model account, while training and inference contribute roughly equally for the recommendation models. Training FLOPs alone omit serving costs, but the larger phase must be established from the service's measured workload and observation period.

Inference workloads, like other software services, scale with demand. An always-available service consumes energy both while processing requests and while maintaining capacity for expected latency and availability. The sustainability account must include query volume, workload mix, and idle provisioned resources.

A production measurement shows how large that account is. Google instrumented the fleet serving its Gemini assistant and attributed energy to the median text prompt in May 2025 across every component of the serving stack (Elsworth et al. 2025). The full-stack figure is 0.24 Wh per prompt (figure 15.12), of which the active accelerators account for only 58 percent. The host processors and memory that feed them add 25 percent, and idle machines held for latency and availability together with facility overhead add another 17 percent. Counting only active accelerators on the most efficiently used machines, the boundary many published estimates use, yields 0.10 Wh, a 2.4× understatement of the same workload. That figure falls below even the accelerator share of the full-stack total because the full stack averages over every serving machine, including those running at lower utilization, while the narrow boundary keeps only the best-utilized ones. The measured prompt also sits at about a quarter of the 1 Wh per query assumed in the worked example above, a reminder that per-query energy is a property of a specific model, serving stack, and utilization

rather than a constant. The same study reports that software and utilization improvements cut the median prompt's energy 33× in a single year, so a per-query figure also carries a date.

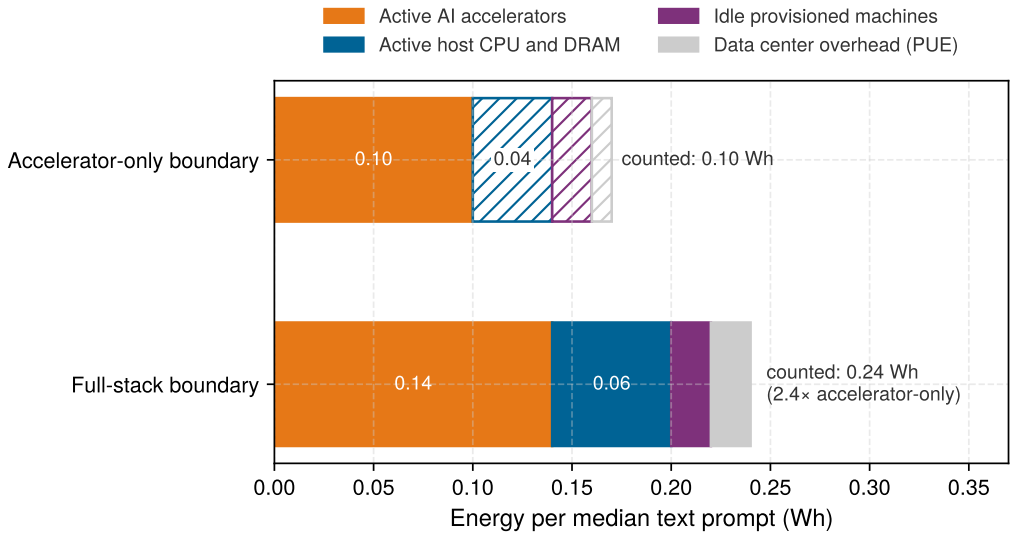

Figure 15.12: Energy per Prompt Depends on the Measurement Boundary: Energy attributed to the median Gemini text prompt in May 2025 under two boundaries. The accelerator-only boundary counts active accelerators on the most efficiently used machines; the full-stack boundary also counts host processors and memory, idle machines provisioned for latency and availability, and data-center overhead, across the whole fleet. Hatched segments are components that the accelerator-only boundary leaves out. Data from Table 1 of Elsworth et al. (2025).

Checkpoint 15.4: The training-inference flip

Consider a vision model where training requires 2,000 GPU-hours at an average accelerator power draw of 300 W. Once deployed, the model serves 1 million requests per day for 730 days, with each request occupying an accelerator for 50 ms at an average draw of 100 W. Assume requests do not overlap on an accelerator. This active-only account excludes idle capacity, host systems, networking, and facility overhead.

- ☐ **Training energy**: calculate the accelerator energy accounted for during training.
- ☐ **Inference energy**: calculate the active-request accelerator energy subtotal over the stated product lifespan.
- ☐ **Optimization focus**: compute the ratio of these inference and training subtotals and identify which phase merits investigation. Explain why choosing the largest achievable lifecycle reduction also requires the omitted energy terms and the savings available from candidate interventions.

15.7.3.1 The energy inefficiency of the decode phase

Prefill and decode stress distinct memory and arithmetic pathways, so meaningful energy comparisons require a declared serving workload. Prefill amortizes loaded weights across all prompt positions simultaneously, whereas each autoregressive decode iteration reuses weights only across active sequence batch slots to emit one token at a time (figure 15.13). Batch size, context length, precision, and the latency target dictate whether memory bandwidth or arithmetic execution limits realized utilization.

↩ Why small-batch decode is bound by memory bandwidth, and how batching reuses each weight load, is established in section 10.6.1.

That split in dataflow governs per-token energy through data movement and arithmetic intensity. During prefill, parallel matrix operations reuse weights across prompt positions and can achieve high compute utilization when the workload exposes enough parallel work. During decode, each iteration

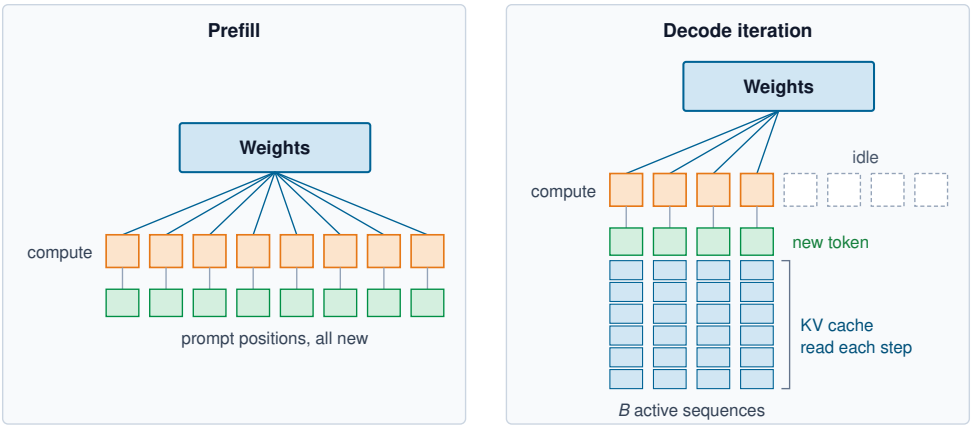

Figure 15.13: Reuse in Prefill and Decode: Prefill reuses weights across prompt positions; a decode iteration can reuse weights across active sequences. At small batches, weight and KV-cache traffic can leave arithmetic units underused. These paths show mechanisms, not measured utilization or a universal energy-per-token ratio.

produces one token per active sequence while accessing model parameters and an expanding KV cache. At small batches, this traffic can leave arithmetic units underused; batching improves weight reuse, subject to memory capacity and latency constraints.

The energy roofline of section 15.4.2.3 puts a number on that underuse. At batch size one, each decode step reads every FP16 weight (2 bytes) to perform two FLOPs with it, an arithmetic intensity of 1 FLOP/byte against an energy crossover of 16 FLOP/byte (equation 15.7). Under the chapter's modeled coefficients, moving each parameter's bytes from DRAM costs 320 pJ while the arithmetic on it costs 20 pJ. Movement is therefore 16× the arithmetic and about 94 percent of the modeled dynamic energy per generated token, before any KV-cache traffic is counted. The coefficient is the table's off-package DRAM value; the high-bandwidth memory that feeds an accelerator costs less per byte (section 2.3.1.1), which narrows the ratio without reversing it at batch one. Batching B sequences reuses each weight load B times and raises the intensity to B FLOP/byte, so the two terms balance only near a batch of 16. Batching is therefore the first energy lever for decode, bounded by memory capacity for the KV cache and by the latency target.

Static power waste occurs when powered circuitry waits rather than producing useful outputs. Reducing that waiting can improve energy efficiency, but a comparison must meter the same hardware boundary and preserve output quality, context, latency, and serving load. Prompt processing and autoregressive generation perform different work; equal token counts alone do not establish a fair energy ratio.

15.7.4 Edge execution and the fleet energy account

Moving inference from a cloud data center onto a billion phones or vehicles does not make its energy cost disappear; it shifts the draw into device batteries, local thermals, and handset manufacturing. Edge intelligence runs models on user devices or nearby computing nodes, avoiding raw-input transfers to a central facility while adding or extending computation at the edge. Local processing and duty cycling reduce total energy only when the communication and remote work they avoid exceed the device's added computation and idle power. Device manufacturing joins that added cost when deployment requires new hardware. Its embodied carbon is allocated over the work the device serves (section 15.5.5), so a lightly used device can carry more manufacturing carbon per inference than the emissions of the transfers that local execution avoids. The comparison holds task quality and service availability fixed.

In autonomous vehicles, time-critical perception and control must meet response deadlines despite variable network connectivity, so onboard computation remains part of the operating budget. Sudhakar et al. (2023) model the potential emissions of a widely deployed autonomous-vehicle fleet,

↩ The power, memory, and compute budgets of edge devices, and how on-device training amplifies them, are established in section 11.3.

showing how device power and operating time can accumulate at global scale; their comparison concerns fleet-level scenarios, not an assertion that each vehicle consumes as much energy as a data center.

Smartphones, wearables, and sensors follow the same accounting. Aggregate operational energy depends on the number of devices, their power states, and the time spent in each state, so a small average draw per device can still produce a large fleet total. Attributing that total to AI requires separating the additional workload from the device's other functions. Device power budgets constrain both local inference and learning, so the comparison must also account for duty cycling, selective cloud escalation, and the repeated computation and model transfers that federated training requires.

15.7.4.1 Edge and mobile power budgets

Edge devices operate within battery, thermal, and supply limits that differ from those of data center accelerators. These limits follow from the deployment rather than the processor's instruction set. The engineering choice is to match each inference workload to a device that meets latency and accuracy requirements within its sustained power and energy budgets.

Power budgets reflect the physical constraints of battery capacity, thermal dissipation, and deployment environment. These constraints propagate differently across the edge AI power budget categories, from microcontrollers to vehicles (table 15.8). TinyML devices (microcontrollers that run small models in kilobytes of memory) must fit average demand to battery-life or harvesting targets while also supporting peak draw. Mobile devices must balance user experience with battery life, and automotive systems face thermal constraints within enclosed vehicle compartments despite having access to vehicle power.

Table 15.8: Edge AI Power Budget Categories: Edge platforms span five orders of magnitude in power consumption, from sub-milliwatt TinyML systems to automotive compute platforms approaching data center power levels. Sustainable deployment requires matching workload requirements to appropriate power tiers.

Platform category	Idle power	Active power	Peak power	Example devices
TinyML (MCU)	1–100 µW	1–50 mW	100 mW	Arduino Nano 33, STM32H7, Nordic nRF5340
Mobile NPU	10–100 mW	0.5–5 W	10 W	Pixel Tensor, Apple Neural Engine, Snapdragon NPU
Edge GPU/TPU	1–5 W	5–30 W	75 W	NVIDIA Jetson Orin NX (10–25 W) and AGX Orin (15–60 W), Google Edge TPU, RPi AI Kit
Autonomous Vehicle	10–50 W	50–200 W	500 W	Tesla FSD Computer, Mobileye EyeQ, NVIDIA Drive

15.7.4.2 TinyML power state dynamics

TinyML devices are the smallest members of edge fleets, sitting below the 2–10 W phones of Archetype C (Federated MobileNet) in both memory and power. Their energy efficiency is a sustainability concern because billions of edge inference calls add up to a measurable carbon footprint at fleet scale. TinyML efficiency depends heavily on duty cycling, where devices alternate between deep sleep and active inference. The duty-cycle model of equation 15.8 applies directly, with $\alpha_{\text{duty}} = t_{\text{inference}}/T_{\text{period}}$ and the deep-sleep draw in place of idle power.

↩ The same duty-cycled power budget, applied to on-device learning rather than inference, is established in section 11.10.1.2.

Consider a keyword-spotting model on a Cortex-M4 microcontroller, an extreme case that obeys the same battery arithmetic as the Archetype C phones. The device draws 15 mW of active inference power for 20 ms in each detection cycle, falls to a deep-sleep draw of 10 microamps at 3.3 V (33 µW) between evaluations, and repeats on a 1 s detection period during continuous listening. Substituting these values into equation 15.8 yields the average power draw:

$$P_{\text{avg}} = 15\ \text{mW} \times 0.02 + 0.033\ \text{mW} \times 0.98 = 0.30 + 0.032 \approx 0.33\ \text{mW}$$

A 250 mAh coin cell at 3 V nominal stores 750 mWh, enough for approximately 2,260 hours, about 94 days, under this idealized budget. The estimate covers compute only. Continuous listening also

requires powered sensing and input buffering while the microcontroller sleeps, and their energy, regulator losses, and battery derating must enter any deployment estimate.

15.7.4.3 On-device learning and the battery wall

Gradient-based on-device learning costs more work and memory than inference alone. A dense gradient step needs roughly 2–3× the forward arithmetic, with activation, gradient, and optimizer-state memory budgeted separately. The battery, thermal, and foreground-latency test that decides when a phone may admit such an update is established in section 11.3.5. Its representative phone envelope (sustained processing within 2–3 W and bursts to 5–10 W for limited periods) bounds how long an on-device update can run alongside battery energy, memory, and the update deadline. Processor TDP alone does not establish feasibility. The sustainability account adds the fleet-scale question of how much energy the admitted updates consume in total.

Napkin Math 15.8: The energy of learning

Problem: Consider fine-tuning a small language model (1B parameters) on a user's smartphone overnight. Is this feasible within a 5 percent battery budget?

Math:

1. Phone battery: Typical capacity is approximately 15 Wh, or about 54,000 J.
2. Budget: 5 percent of 54,000 J = 2,700 J.
3. Training cost: this step uses illustrative per-parameter energy coefficients for a mobile accelerator (assumed for this scenario, not measured on a particular phone). The backward pass costs twice the forward pass because it computes gradients with respect to both activations and weights, so one training step costs about three forward passes:
   - Forward pass: ≈ 2 nJ/param.
   - Backward pass: ≈ 4 nJ/param.
   - Total per token: 6 nJ/param $\times 10^9$ params = 6 J/token.
4. Capacity: 2,700 J / 6 J/token = 450 tokens.

Systems insight: The calculation establishes a token capacity, not whether the proposed adaptation fits. That decision also requires the token count needed to reach the quality target, the available memory, and sustained device power. If the required work exceeds the battery budget, parameter-efficient fine-tuning or sparse updates are candidates for reducing cost; their measured energy and achieved quality must establish whether they close the gap.

The local-processing decision compares the energy of the required computation with the communication and remote computation it avoids. Per-bit radio costs motivate the comparison, but the full job must also include repeated training rounds and model transfers.

Systems Perspective 15.3: The energy hierarchy

Trade-off: The architectural choice between processing data locally and sending it to the cloud is governed by an energy budget. The energy-to-communication ratio motivates comparing local work with transmission, but a decision requires the total energy of each alternative at the required quality and latency.

Energy cost per operation (approximate):

- 32-bit integer add: 0.1 pJ
- 32-bit float mult: 3.7 pJ
- Wireless transmit (1 bit): 100,000–500,000 pJ (Bluetooth/Wi-Fi)

Systems insight: Under these illustrative component costs, a transmitted bit consumes far more energy than one arithmetic operation. The radio cost is quoted per bit; per byte it reaches

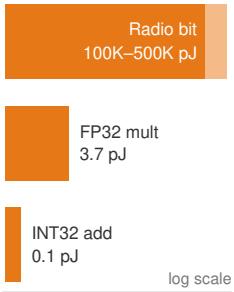

Radio transmission dwarfs local arithmetic in energy.

roughly 0.8 to 4 million pJ, about two orders of magnitude above the lower bound given for the wired network row of table 15.2. This motivates local processing but does not establish an offloading threshold: include memory, radio startup and protocol overhead, idle power, and remote computation in the full edge-cloud comparison. Local processing wins only when those complete terms favor it.

15.7.4.4 Energy harvesting for autonomous edge AI

TinyML can operate from harvested ambient energy when the available supply supports the full device workload. Sustained operation requires both an adequate long-term energy budget and a way to handle intervals when harvesting falls below demand, through storage, duty cycling, or suspended execution.

Ambient energy harvesting delivers from about a microwatt to about ten milliwatts per device, depending on the source (table 15.9). A keyword spotting model optimized to 0.5 mW average power matches the nominal output of approximately 5 cm² of indoor solar harvesting only near the top of the listed illumination range, before conversion losses and other device loads. This average-power comparison does not establish uninterrupted operation. The design must cover sensing, storage losses, and periods of dim or absent light, or permit the workload to pause. Ambient supply removes the need for purchased operating electricity only when it covers that complete budget; manufacturing, storage, maintenance, and disposal still enter the lifecycle account.

Table 15.9: Energy Harvesting Power Budgets: Ambient energy harvesting enables batteryless TinyML deployments when average power consumption remains within harvesting capacity. Solar harvesting provides the highest power density for most deployments.

Harvesting source	Typical power	Viable TinyML applications
Indoor solar (1 cm²)	10–100 microwatts	Periodic sensor classification
Outdoor solar (1 cm²)	1–10 milliwatts	Continuous keyword spotting
Thermoelectric (body heat)	10–100 microwatts	Wearable gesture recognition
RF harvesting (Wi-Fi)	1–10 microwatts	Ultra-low-duty sensor nodes
Vibration piezoelectric	100 microwatts–1 mW	Industrial monitoring

15.7.4.5 Cascade inference architecture

Beyond individual device efficiency, architectural patterns determine total system energy consumption across edge-cloud boundaries. A cascade architecture deploys a task-appropriate small edge model to filter inputs before cloud inference. Equation 15.14 expresses total energy as the sum of local processing plus probabilistically-triggered cloud costs:

$$E_{\text{cascade}} = E_{\text{local}} + p_{\text{escalate}} \times (E_{\text{transmit}} + E_{\text{cloud}}) \tag{15.14}$$

where p_{escalate} is the probability of requiring cloud inference (measured on the deployment input distribution).

Consider an illustrative visual inspection system where a lightweight edge model, such as MobileNetV3-Small, consumes 0.5 mJ per image classification, whereas a high-capacity cloud model, such as ResNet-152, requires 50 mJ per inference. Transmitting an image across a cellular uplink adds 10 mJ. If the edge model resolves routine cases locally and escalates only ambiguous inputs, for an escalation rate p_{escalate} of 10 percent, substituting these parameters into equation 15.14 gives the average cascade energy:

$$E_{\text{cascade}} = 0.5 + 0.10 \times (10 + 50) = 0.5 + 6.0 = 6.5 \text{ mJ/image}$$

Compared to always-cloud inference at 60 mJ per image, this illustrative cascade reduces the modeled energy by 89 percent. Quality requires a separate test of local decisions, missed escalations, and the cloud model; the energy calculation says nothing about accuracy.

15.7.4.6 Wake-word triggered systems

Always-on systems can use hierarchical wake detection to reduce average power across stages of increasing capability and draw. In this illustrative model, an analog front end consumes 10 µW while monitoring for voice activity. A detected event activates a compact neural network wake detector drawing 100 µW. A positive wake decision then triggers full-model inference at 10 mW for 50 ms. The energy calculation uses those assumed state powers and event rates; a deployed system must also account for false activations, missed wake words, and transition costs when evaluating its energy and service quality.

For an illustrative speech-active fraction of 5 percent and a wake-event rate of 0.001 per second (one event per 1,000 s), each full-model event lasts 0.05 s. The front end draws continuously, the wake detector draws only during speech, and the full model contributes its event rate times its event energy:

↩ The always-on budget of wake-word detection, set by how often the detector evaluates over the listening period, is established in section 11.2.2.

$$P_{\text{avg}} = 0.01 + 0.05 \times 0.1 + 0.001 \times 10 \times 0.05 = 0.0155 \text{ mW}$$

The hierarchical approach achieves 15.5 µW average power compared to 10 mW for always-active full inference, a reduction of roughly 645× within this modeled boundary. Battery lifetime also requires battery capacity, conversion losses, sensors, and other device loads.

15.7.4.7 Federated learning energy analysis

Training at the edge keeps raw data local but replaces its transfer with repeated model downloads and client-update uploads, one of each per participating client in every round of federated averaging. This is the regime of Archetype C (Federated MobileNet), where phones train a shared model within a sustained power budget of a few watts. Equation 15.15 contrasts the energy trade-offs between federated and centralized approaches. Here R is the number of rounds, and N_{clients} is the number of participating clients per round. The term E_{upload} is one client-update upload, E_{raw} is one client's raw-data upload, and $E_{\text{cloud,train}}$ is the energy of one centralized cloud training run:

$$E_{\text{federated}} = R \left[N_{\text{clients}} \left(E_{\text{client}} + E_{\text{download}} + E_{\text{upload}} \right) + E_{\text{aggregation}} \right]$$
$$E_{\text{centralized}} = N_{\text{clients}} \times E_{\text{raw}} + E_{\text{cloud,train}} \tag{15.15}$$

↩ Federated averaging, in which each round sends the model to participating clients and aggregates their updates, is established in section 11.6.2.2.

Federated learning becomes more energy-efficient only when repeated model communication, local training, and aggregation consume less energy than raw-data transfer and centralized training at comparable model quality. The comparison must cover the same client population and account for repeated participation across rounds. Data being larger than one model update is insufficient by itself because downloads and uploads recur. Privacy *and* energy benefits require separate evidence: keeping raw data local does not by itself prevent updates from revealing private information, and any protection mechanism adds its own costs to the energy account.

Training, serving, and edge deployment all draw operational energy, and none of these accounts includes the hardware that performs the work. That physical supply chain, from fabrication through disposal, is the lifecycle term that remains to be counted.

15.8 Hardware Lifecycle and E-Waste

The environmental cost of an AI accelerator begins long before its first FLOP is calculated (Gupta et al. 2022; Luccioni et al. 2023; NVIDIA Corporation 2025; Schneider et al. 2025). The per-H100 manufacturing footprint quantified in section 15.5.5 is incurred before the device is ever powered on.[17] A training fleet of a few thousand such accelerators therefore carries a manufacturing debt of the same order as the operational emissions of a GPT-3-scale training run on the US-average grid (about 552,123 kg).

A comprehensive **life cycle assessment (LCA)** quantifies the cumulative environmental impact across four key phases: design, manufacture, use, and disposal. LCA can reveal manufacturing as a major share of lifecycle impact, making it a critical sustainability lever that operational efficiency improvements alone cannot address. The first published LCA of an AI accelerator, covering five generations of Google's TPU across all four phases, introduced compute carbon intensity, the lifecycle emissions per unit of computation. Unlike performance per watt, the metric charges each chip for its

17 **Life cycle assessment (LCA)**: Standardized by ISO 14040/14044, LCA traces environmental impact from raw material extraction through disposal (International Organization for Standardization 2006a, 2006b). For AI hardware, LCA separates embodied manufacturing emissions from operational energy use, making hardware refresh cycles and accelerator lifespan extension first-order sustainability levers that operational efficiency alone cannot substitute.

own manufacturing, and by that measure the TPU v6e is about 3× better than the TPU v4i (Schneider et al. 2025).

These four phases organize the same emissions by time rather than by activity or reporting scope, as the second column of table 15.5 records, so a phase total must never be added to an activity or scope total. Evaluated across that full timeline, life cycle assessments can show that discarding functional hardware purely for modest efficiency gains causes more environmental harm through embodied carbon than it saves in operational power. The tipping point where new hardware becomes environmentally justified depends on utilization and grid intensity, a break-even worked in section 15.8.6.

Each of the four lifecycle stages adds its burden through a distinct physical mechanism (figure 15.14). System design fixes much of the lifetime load, manufacturing embeds carbon, water, and chemical burdens during fabrication, operation draws grid electricity for training and serving, and decommissioning produces electronic waste unless hardware is refurbished or its components recovered.

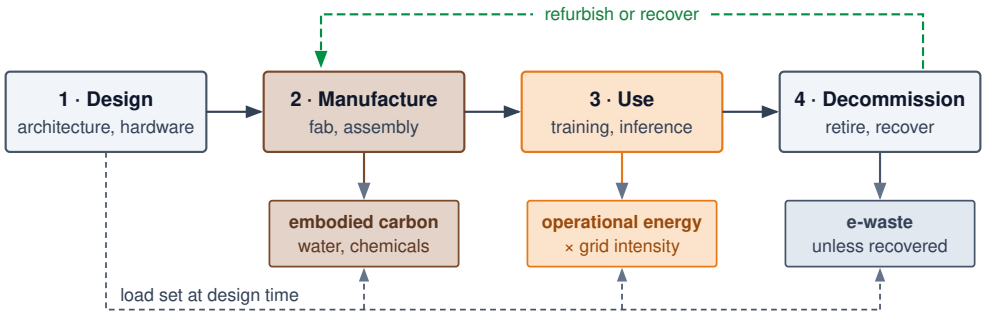

Figure 15.14: AI System Lifecycle: Four stages in sequence, each dropping the burden it adds, with a recovery loop from decommissioning back to manufacture. Design decisions fix much of the load the later stages carry; manufacturing adds embodied carbon, water, and chemicals; use adds operational energy at the grid's carbon intensity; decommissioning adds e-waste unless hardware is refurbished or its components recovered.

The lifecycle sequence connects design choices to costs incurred at different times. Experimentation consumes resources before deployment; manufacturing adds embodied carbon; and use couples the workload to grid and cooling constraints. Retirement creates a choice among reuse, recovery, and disposal. Extending useful service spreads the hardware's embodied burden over more work, but it does not reverse emissions already incurred. The environmental benefit depends on the replacement it avoids and the operating cost of keeping the existing equipment in service.

15.8.1 Design and experimentation phase

Public training-energy figures almost always report the single final run that produced the published checkpoint, omitting the months of failed experiments, hyperparameter sweeps, and architecture searches that preceded it. Those preliminary runs belong to the design phase, which covers research, development, and optimization before deployment. Automated neural architecture search (NAS) can evaluate many configurations. Section 15.7.1 showed that two estimates of the same search can differ by orders of magnitude depending on whether the boundary includes proxy searches, full runs, data-center overhead, grid mix, and reuse across downstream models (Strubell et al. 2019; Patterson et al. 2021). Later efficient NAS methods use weight sharing, continuous relaxations, and hardware-aware search to reduce that search budget (Elsken et al. 2019). Excluding search does not make the remaining figures comparable either, because final runs also resist ranking by model size (table 15.10).

Table 15.10: Final-Run Training Footprints: Reported energy and operational emissions of five large training runs (Patterson et al. 2021). Each row is a final run under its authors' boundary; search, failed runs, and evaluation are excluded, and the grid factor differs by data center. Compare footprints only after aligning workload, energy boundary, grid factor, and development scope.

Model	Parameters	Energy (MWh)	Emissions (t CO_2e)	t CO_2e per MWh
T5	11B	86	47	0.55
Meena	2.6B	232	96	0.41
GShard-600B	619B	24	4.3	0.18
Switch Transformer	1.5T	179	59	0.33
GPT-3	175B	1,287	552	0.43

The reported final run for the 600-billion-parameter GShard model emitted under a tenth as much as the 11-billion-parameter T5 run. One reason is architectural, because GShard-600B and the 1.5-trillion-parameter Switch Transformer are Mixture-of-Experts models that route each token to only a small subset of experts rather than activating every parameter on every token as dense T5, Meena, and GPT-3 do. The table also records both lower energy use and lower emissions per megawatt-hour for GShard. Emissions per megawatt-hour is each run's effective grid factor, and across the five runs it varies by a factor of about 3. Because the runs also differ in token budgets, accelerator generations, and data-center grids, the table's totals do not isolate the saving from sparse activation alone. Parameter count therefore cannot rank their emissions, and operation count would still need an execution-efficiency and grid account. Search and failed experiments require a separate development inventory before these final-run figures can support a lifecycle comparison.

Addressing the design phase's sustainability challenges requires innovations in training efficiency: sparse training, low-precision arithmetic, weight-sharing, and energy-aware NAS approaches. Transfer learning and fine-tuning pretrained models reuse pretrained representations instead of requiring every task to be trained from scratch (Raffel et al. 2020).

↩ Mixture-of-experts execution and the AllToAll routing it requires are established in section 6.5.5.

15.8.2 Water, chemicals, and critical materials

Before an accelerator exists, its fabrication consumes physical resources that an energy account alone does not quantify: ultrapure water, process chemicals, and materials. Their burdens occur across extraction, refining, fabrication, and assembly sites, rather than solely where the finished hardware operates. Water demand must be evaluated against local availability and competing uses, while material supply depends on the component and manufacturing process (Jha 2014; Chen 2006). These resource accounts extend the lifecycle boundary beyond carbon without reducing unlike impacts to a single energy total.

Water accounting requires more than a single WUE ratio. **Water withdrawal** measures water removed from a source, whereas **water consumption** measures the portion not returned to the same watershed in a usable form. The account must separate direct on-site cooling water, which equation 15.3 and the cooling-tower calculation in section 15.6.3 estimate, from indirect water used to generate electricity (figure 15.10) and manufacture hardware. The account must then attach each flow to its watershed, season, and local scarcity (Environmental Management—Water Footprint—Principles, Requirements and Guidelines 2014; The Green Grid 2011). One liter consumed during a wet season in a low-stress basin is not equivalent to one liter consumed during a drought in a stressed basin, so decision-grade comparisons report both physical volume and a stress-weighted result. Reclaimed-water percentages also need a declared denominator because recirculating water inside a facility can reduce withdrawal without eliminating evaporative consumption.

Semiconductor fabrication uses ultrapure water for chemical-mechanical planarization, wet etching, and wafer rinsing (Cooper et al. 2011).[18] Complete environmental accounting requires auditing gross process flow, freshwater withdrawal, recycled volume, and evaporative consumption as distinct quantities, because internal reuse and flows that leave the facility change different totals (figure 15.15). Treatment returns reclaimed water to the intake stream and can reduce freshwater withdrawal,

[18] **Semiconductor water scale**: A fab water inventory must separate intake, internal reuse, discharge, and consumption. Water per usable die depends on process flow, yield, allocation, and site conditions; a smaller process-node label does not establish a universal per-die multiplier.

whereas evaporation remains a consumptive loss within the local water account. That classification does not mean the water permanently disappears from the hydrological cycle.

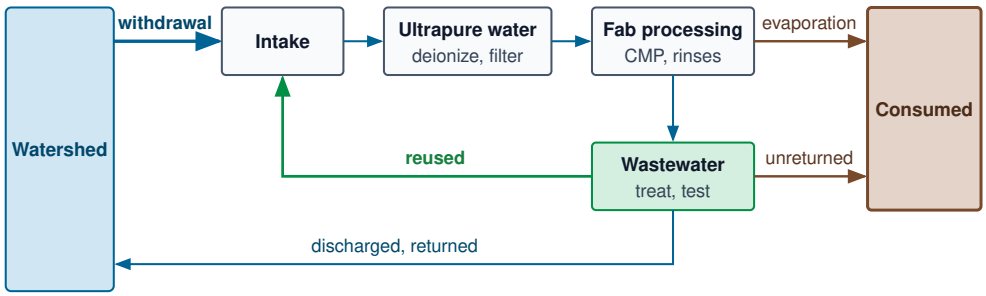

Figure 15.15: Semiconductor Water Cycle: Intake supplies process water; treatment returns part of the flow for reuse. Recycling can reduce freshwater withdrawal, while evaporation and other unreturned flows determine consumption. State the denominator of any reuse fraction.

The separate flows in figure 15.15 show why a recycling fraction cannot determine absolute consumption. The same reuse percentage can accompany very different intake volumes and evaporative losses. Facility siting therefore requires the underlying water volumes and local watershed conditions, not a reuse percentage alone.

Fabrication also uses chemical products for etching, doping, and cleaning: strong acids such as hydrofluoric and sulfuric acid, organic solvents, and toxic dopant gases such as arsine and phosphine. Their hazard comes from toxicity, reactivity, and the treatment burden of the waste streams they produce, not from their energy content. Chemical inventories and safety data sheets identify hazardous ingredients, although undisclosed formulations can limit the assessment (S. Kim et al. 2018). The environmental account needs the substances and quantities used at the specific facility, together with their release and treatment pathways. An energy total alone cannot describe these chemical burdens.

AI hardware and its manufacturing supply chain depend on **critical materials** whose supply disruption can constrain production. Criticality is not synonymous with geological rarity or membership in the rare earth group. Materials such as indium appear in critical-materials and endangered-elements analyses because supply can depend on byproduct extraction, substitution options, and recycling constraints (Rhodes 2019). The relevant dependency may lie in a processor, package, peripheral component, or fabrication process (table 15.11).

Table 15.11: Critical Materials for AI Hardware: Supply risk comes from different sources for different materials, such as byproduct extraction, geographic concentration, or conflict sourcing, rather than from geological rarity alone. A dependency on any one of them can constrain hardware production even when energy and carbon budgets are met.

Material	Application in AI semiconductor manufacturing	Supply concerns
Silicon (Si)	Primary substrate for chips, wafers, transistors	• Processing constraints • Geopolitical risks
Gallium (Ga)	GaN-based power amplifiers, high-frequency components	• Limited availability • Byproduct of aluminum and zinc production
Germanium (Ge)	High-speed transistors, photodetectors, optical interconnects	• Scarcity • Geographically concentrated
Indium (In)	Indium tin oxide, optoelectronics	• Limited reserves • Recycling dependency • Conflict mineral • Vulnerable supply chains • High geopolitical risks • Environmental extraction concerns
Tantalum (Ta)	Capacitors, stable integrated components	
Rare earth elements	Magnets, sensors, high-performance electronics	

Material	Application in AI semiconductor manufacturing	Supply concerns
Cobalt (Co)	Batteries for edge computing devices	• Human rights issues • Geographical concentration (Congo)
Tungsten (W)	Contact and via plugs, diffusion barriers	• Limited production sites • Geopolitical concerns
Copper (Cu)	On-chip interconnect wiring, heat spreaders	• Limited high-purity sources • Geopolitical concerns
Helium (He)	Semiconductor cooling, plasma etching, EUV lithography	• Nonrenewable • Irretrievable atmospheric loss • Limited extraction capacity

The construction and operation of fabs and data centers also directly impact natural ecosystems through habitat disruption, water stress, and pollution from chemical discharge. Semiconductor-industry effluents can contaminate nearby fluvial sediments with heavy metals and trace elements (Hsu et al. 2016). Waste generation from fabrication, including gaseous emissions, VOC-laden air, and metal-contaminated wastewater, requires advanced treatment systems, and end-of-life hardware adds the e-waste stream examined in section 15.8.5.

15.8.3 Manufacturing phase

Fabricating an AI accelerator near the reticle limit, the maximum die area a single lithography exposure can print, leaves far fewer candidate dies per wafer than a small mobile processor, and every failed die transfers its share of ultrapure water, process chemicals, and lithography energy onto the working ones that survive. The manufacturing footprint combines the embodied carbon established in section 15.5.5 with the water, chemical, and material demands in section 15.8.2, and two structural properties of modern accelerators enlarge that resource investment per usable device.

First, because the probability that a die escapes every random wafer defect falls roughly exponentially with die area, a defect density that kills 5 percent of small dies may kill 20 percent or more of reticle-limit dies, concentrating the wafer's fabrication burden onto a smaller surviving yield. Second, the memory bandwidth that large-scale training and inference require pushes accelerators toward advanced 2.5D and 3D packaging, such as TSMC's Chip-on-Wafer-on-Substrate (CoWoS), which places high bandwidth memory stacks beside the logic die on an interposer. The inventory must include those memory stacks, the interposer, and the added assembly stages rather than counting only the accelerator die. Each usable accelerator therefore carries more process water, chemicals, and energy than its logic-die area alone suggests. Longer useful service spreads that manufacturing investment across more work, while lower-toxicity materials, recycling, and cleaner manufacturing address the production burden itself (Cenci et al. 2021; Irimia-Vladu 2014).

15.8.4 Use phase

A GPT-3-scale training run draws 1,287 MWh over weeks of sustained cluster operation, and when those megawatt-hours are drawn matters just as much to the grid and the atmosphere as their sum. The operational energy consumed during training and inference is detailed in section 15.7; how flexible workloads shift that draw toward cleaner hours is developed in section 15.9.4. Once that active service ends, the remaining lifecycle question is what happens to the physical equipment itself.

15.8.5 Disposal, e-waste, and embedded AI

Data centers routinely retire working AI servers after three to four years, not because the silicon has physically failed, but because a newer generation's capacity or efficiency leaves older racks uncompetitive for frontier workloads. Those hardware upgrade cycles (Slade 2007) connect AI deployment to the broader electronic waste (e-waste) problem. In 2019, global e-waste reached roughly 54 million metric tons (Forti et al. 2020), of which 17.4 percent was documented as formally collected and recycled (Singh and Ogunseitan 2022). These totals cover electrical and electronic equipment generally, not AI hardware alone. When equipment is finally discarded, hazardous substances in its components require controlled recovery and treatment rather than uncontrolled dumping or processing (Grossman 2007).

Before hardware ever reaches that waste stream, memory capacity, bandwidth, and interconnect compatibility govern whether an older rack can be upgraded in place or must be retired as a unit. In tightly integrated systems, a new accelerator generation may require baseboard, networking, or server changes, whereas PCIe-based designs can offer modular upgrade options. Evaluating the actual compatibility boundary separates commercial retirement from physical disposal, because a functioning node retired from frontier training can still serve less demanding batch, development, or inference workloads.

The deployment of **embedded AI** integrates inference into edge appliances and consumer electronics. The broader connected-device fleet (figure 15.16) grew from an estimated 8.6 billion devices in 2019 to 15.1 billion in 2023, with 29.4 billion projected by 2030 (Statista 2022). These counts indicate the scale of the device ecosystem, not the number of AI systems or the amount of e-waste. Waste generation also depends on service life, repair, and reuse. Adhesive enclosures, nonreplaceable batteries, and proprietary components can make repair difficult, increasing the risk that a limited failure leads to whole-device disposal (Baldé et al. 2017).

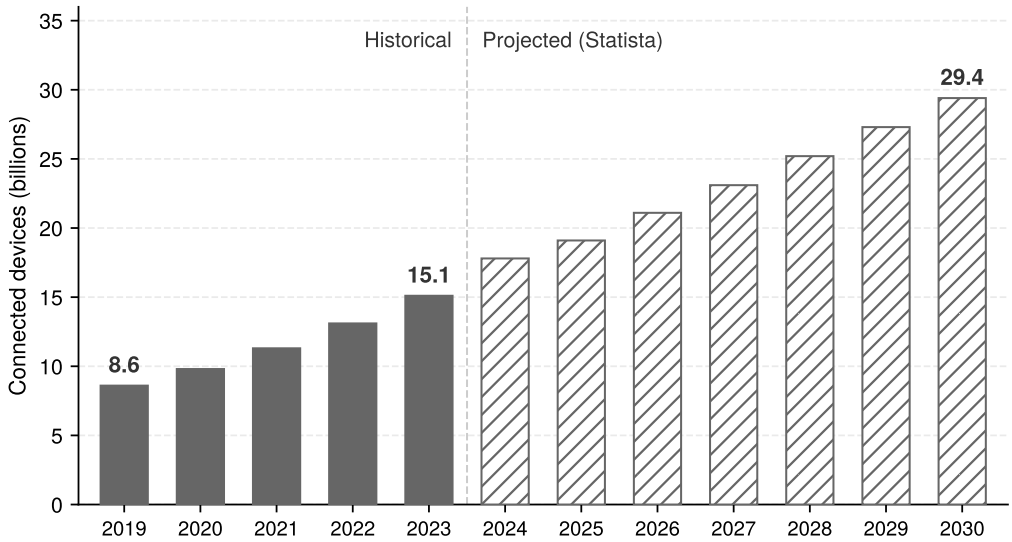

Figure 15.16: IoT Device Growth: Connected devices worldwide, with historical counts through 2023 (unhatched) and projections from 2024 onward (hatched) (Statista 2022). The counts describe the connected-device ecosystem, not the number of AI systems or the amount of e-waste. Waste generation also depends on service lifetimes, repair, reuse, and disposal.

Limited software support, restricted repair, and difficult disassembly can push devices toward replacement instead of reuse. The lifecycle boundary must then follow the retired equipment to its next use or treatment destination. Moving hardware out of the operator's inventory does not remove its environmental burden; the account must distinguish useful reuse, controlled recovery, and disposal that transfers costs to other communities.

15.8.6 Hardware lifespan extension

Keeping an AI server in useful service for 5 years instead of 3 years cuts the embodied carbon allocated to each year of operation by 40 percent. Countering the linear "take-make-dispose" model requires a shift toward a **circular economy** (Stahel 2016) that prioritizes reuse, refurbishment, and recycling, spreading a fixed manufacturing footprint over more service. This allocation result motivates evaluating longer use, but a claim of avoided emissions requires comparing the replacement and operating costs of the alternatives.

Four lifecycle interventions support longer hardware service by making systems repairable, upgradeable, supported, and reusable across their physical life. Physically, right-to-repair initiatives reduce service restrictions, including barriers to replacement parts, tools, and repair information (Federal Trade Commission 2021). Modular design separates replaceable subsystems so that an

accelerator or network upgrade need not require discarding the entire chassis. Once the hardware remains serviceable, extended support cycles provide the driver updates and security patches needed to keep functioning equipment usable. When a node eventually falls behind frontier training demands, secondary-use programs match retired accelerators to workloads they can still serve, such as batch processing, development, or research. Each intervention must preserve the performance, compatibility, and security required for useful continued operation. These interventions move sustainability from disposal management to lifecycle engineering.

A replacement decision compares two ways of delivering the same future service, continued operation of the existing fleet and operation of a replacement fleet. The original manufacturing emissions have already occurred under either choice. Replacement adds a new manufacturing burden that pays back only through operational savings over the remaining service interval, so the tipping point is the time those savings take to equal it.

Suppose a replacement accelerator carries the same allocated manufacturing footprint as an H100 (164 kg CO_2e) and completes the same work with 50 percent less energy. Against an existing accelerator running continuously at the rated 700 W, the replacement repays its manufacturing emissions in about 6.5 weeks on the US-average grid. On Quebec's grid it repays them only after about 2.7 years, nearly the whole of a 3-year standard service life. Lower utilization stretches both intervals in proportion. The same upgrade can therefore repay itself within weeks on a fossil-heavy grid and become a net loss on a clean one if the replacement leaves service before it breaks even. This accelerator-only account omits host and facility overhead, and disposal and any useful second life for the retired hardware shift it further, so a lower annual amortized figure alone cannot settle the decision (the financial side of the same refresh decision is screened in section 12.9.3.1).

Once an engineering team has measured both the operational and lifecycle terms of its workload account, the remaining question is how to combine algorithmic, infrastructure, and scheduling levers to drive those terms down under a real power or carbon ceiling.

15.9 Mitigation Strategies

When a data center hits its absolute power ceiling, the operator cannot simply buy more GPUs. Supporting additional useful work requires reducing the power needed for that work, changing its schedule, or securing additional feasible capacity. Algorithmic optimization can reduce executed work and data movement; infrastructure changes can reduce overhead. Carbon-aware placement addresses a different constraint by changing the emissions associated with electricity use, even when energy demand stays fixed.

The measurement frameworks identify where environmental costs concentrate in a particular workload. Repeated experiments can make development dominant; sustained serving can make inference dominant; and clean electricity or limited use can increase the relative importance of manufacturing. The measured balance guides intervention. Algorithmic changes can reduce the number or cost of operations, infrastructure changes affect execution and facility overhead, and governance sets the budgets and incentives under which those choices operate.

Implementation must account for the rebound introduced with fleet growth in section 15.3, known as the Jevons paradox[19] (principle 25). The tenfold-efficiency, hundredfold-volume example there supplied the arithmetic; the mechanism is that cheaper computation enables applications that were previously economically infeasible, so usage can grow enough to erase or exceed the expected energy savings. This **rebound effect** makes absolute outcomes, including carbon budgets, necessary companions to rate efficiency (FLOP/s per watt). Renewable sourcing changes the emissions associated with electricity but does not cap consumption. Governance must track workload growth against the intended resource budget to determine whether efficiency gains reduce the fleet's total impact.

[19] **Jevons paradox**: Jevons (1865) described how efficiency could stimulate demand enough to increase total coal use. Price reductions alone do not establish either energy factor; measure both unit energy and demand.

15.9.1 Multi-layer mitigation

Compressing a model in half saves no carbon if the freed cluster capacity immediately fills with new experiments or if the job runs on a coal-heavy grid without a facility cap. Because a local efficiency gain can itself raise total consumption, mitigation must decide both which layer owns each reduction and which absolute budget contains the growth in demand that follows. Energy-efficient model design, optimized hardware deployment, sustainable infrastructure operations, and carbon-aware

scheduling each attack a different term in the lifecycle footprint. Framework selection matters only insofar as it changes computation, memory movement, utilization, or reporting. Lifecycle-aware design keeps the optimization from ending at deployment by checking whether savings survive training, inference, manufacturing, and use.

Each mitigation lever acts at one accounting tier (figure 15.17). Work reduction and execution efficiency change IT energy. Facility overhead scales that energy to the facility total. A regional grid factor converts facility kilowatt-hours into operational carbon. Hardware amortization adds embodied emissions at its own boundary. Reductions at the multiplicative tiers compound only when each is measured in the combined configuration on the same fixed useful work, because a cut in IT energy can itself change PUE. Isolated efficiency percentages measured on different workloads or baselines, or applied to the additive embodied term, cannot be multiplied into a single composite reduction factor.

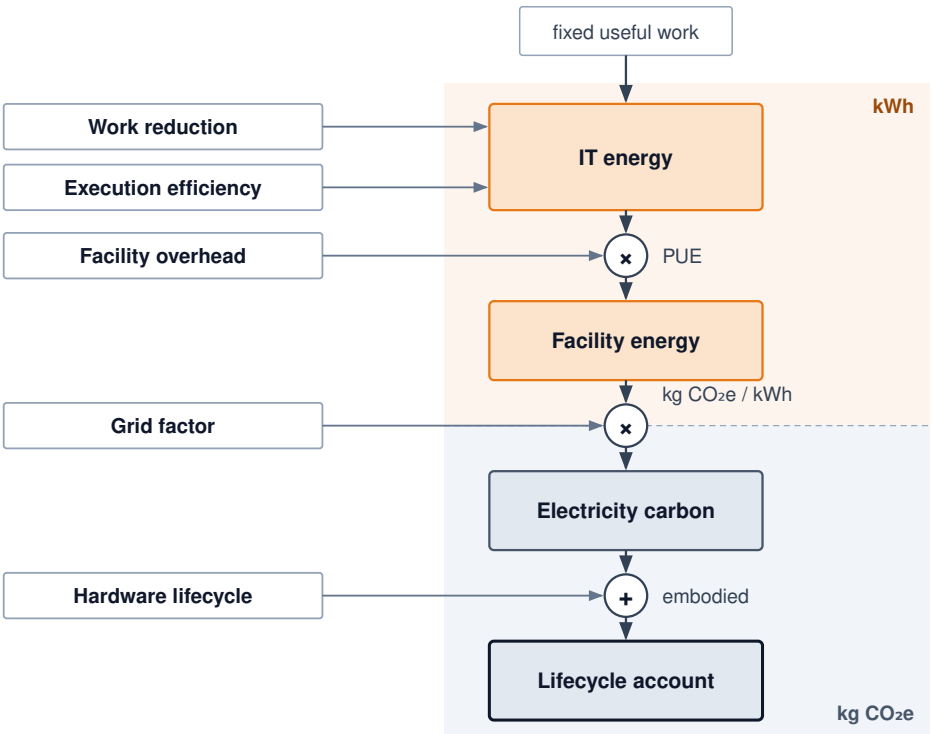

Figure 15.17: Interventions and Their Accounting Boundaries: Work and execution changes affect IT energy; facility overhead converts that energy to a facility total. A chosen grid factor converts electricity to emissions, while hardware lifecycle emissions enter separately. Evaluate the complete configuration at fixed useful work rather than multiplying unrelated improvement factors.

Rebound can cancel a reduction at any of these tiers. On an elastic demand curve ($\epsilon_{\text{elast}} > 1$), where volume changes by a larger fraction than unit cost, a fall in unit cost raises volume enough to outweigh the saving (figure 15.18). When unit cost falls from P_0 to P_1, the green shaded area represents the saving on baseline volume Q_0. Demand expands to Q_1, creating a red area ($P_1(Q_1 - Q_0)$) that exceeds the original saving in this schematic. With monetary unit cost on the vertical axis, the areas measure expenditure. They represent energy only under an additional assumption that unit cost changes in proportion to energy per operation; an emissions comparison also needs the applicable carbon intensity.

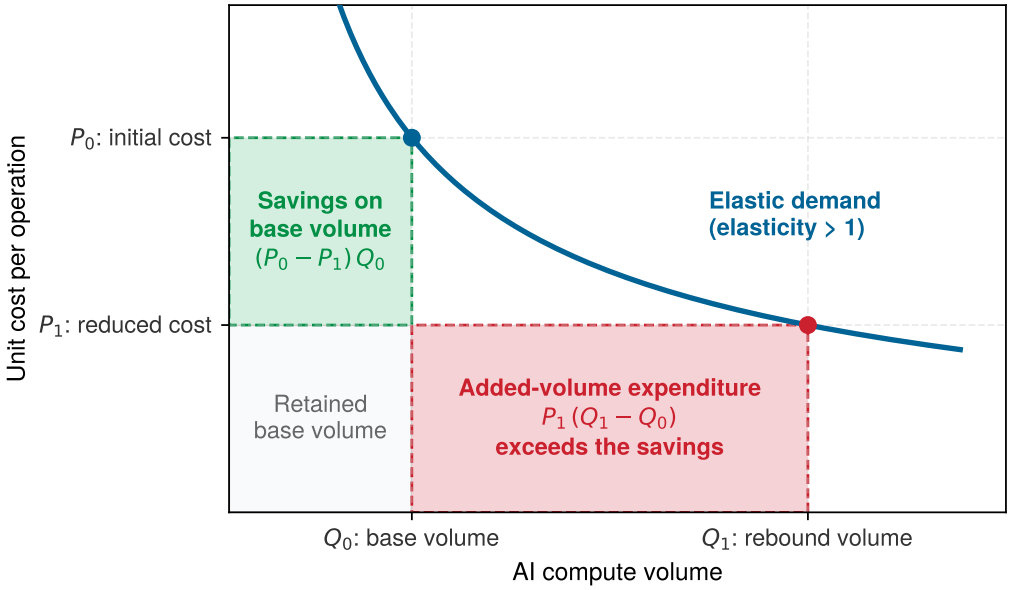

Figure 15.18: Jevons Paradox in AI Workloads: A schematic elastic demand curve. When unit cost falls from P_0 to P_1, demand grows from Q_0 to Q_1. The green rectangle is the expenditure saved on the original volume; the red rectangle is the expenditure on added volume at the lower unit cost. In this schematic, the red area exceeds the green. These areas represent energy only if unit cost changes proportionally to energy per operation; an emissions comparison also requires carbon intensity.

Total energy depends on measured energy per query and query volume. Demand elasticity describes how volume responds to price; translating that response into energy requires a separate account of how unit energy changes.

Checkpoint 15.5: The efficiency trap (Jevons paradox)

An engineering team optimizes a translation service, reducing measured energy per query by 50 percent (2× efficiency gain) at fixed quality. Assume the service passes that saving through as a proportional price reduction and that grid intensity stays unchanged.

- ☐ **Inelastic demand**: calculate the change in total energy consumption when demand is inelastic and price change does not affect usage.
- ☐ **Elastic demand**: calculate the net change in total energy consumption when a 50 percent cost reduction leads to a 300 percent increase in query volume.
- ☐ **Rebound effect**: define how this "rebound effect" challenges the assumption that "efficient models are automatically green models."

The Jevons paradox does not invalidate efficiency as a strategy; it makes governance and capacity planning part of the result. At the level of individual systems, reducing energy per useful task remains a direct sustainability lever. A reduction in latency or monetary cost is a candidate improvement whose energy effect must still be measured.

Systems Perspective 15.4: Efficiency as sustainability

Model optimizations become sustainability tools only when they reduce measured lifecycle cost at fixed quality. Pruning can remove computation and memory traffic when the hardware skips the resulting structure. Quantization reduces representation size and may accelerate

supported kernels, while knowledge distillation trades additional training for a smaller serving model. Each technique can also fail to reduce wall energy when its overhead or deployment regime leaves the dominant term unchanged.

Performance engineering supplies candidate interventions, but the environmental result still needs measurement. A faster run reduces energy only if the shorter duration outweighs any increase in average power; lower memory use helps only if it removes powered work or capacity. Hardware-software co-design becomes a sustainability gain when these savings survive the lifecycle boundary and the deployed workload volume.

Sustainable AI engineering overlaps strongly with efficient AI engineering but extends beyond it. The engineering principles that enable systems to scale, perform better, and cost less to operate can also reduce environmental impact, but sustainability adds lifecycle accounting, carbon-aware placement, absolute resource budgets, water and materials constraints, and governance against rebound effects. Sustainability is an integral part of good systems engineering, not a synonym for efficiency alone.

15.9.2 Lifecycle-aware development

Spending weeks quantizing a model whose footprint is dominated by repeated architecture searches, or moving a training job to a cleaner grid when year-round serving draws ten times more electricity, solves the wrong part of the inventory. Lifecycle-aware development starts from the largest environmental term in the workload and then selects the intervention that changes that term. Algorithmic design, infrastructure optimization, operational practice, and governance reduce impact only when they are sequenced around the measured bottleneck rather than applied as a generic checklist (Uddin and Rahman 2012).

15.9.2.1 Energy-efficient algorithmic design

Large deep learning models can require substantial computation during training and inference. The engineering question is how much of that work a particular task needs at its required quality. Parameter count alone does not identify redundant computation, and removing weights saves energy only if the resulting implementation avoids work or data movement. Sustainable algorithm design therefore evaluates the model and its execution together, measuring energy per useful task rather than treating a smaller representation as sufficient evidence of improvement.

When the measured bottleneck is unused structure, bit width, or serving scale, the sustainable design lever changes, because each bottleneck has its own intervention on the energy term (table 15.12).

Table 15.12: Algorithmic Energy Levers by Bottleneck: Pruning, quantization, and distillation reduce energy through different mechanisms, so the right choice depends on the measured bottleneck rather than on a generic compression goal.

Technique	Measured bottleneck	Energy mechanism	Representative evidence
Pruning	Unused model structure	Removes redundant weights, reducing model size, compute, and memory movement during inference	Structured pruning can remove complete channels or blocks; measure energy and quality for the resulting executable model
Quantization	Bit width	Lowers numerical precision so arithmetic units and memory transfers move fewer bits	Q8BERT demonstrates an 8-bit BERT representation with a smaller model footprint and task-dependent quality retention (Zafrir et al. 2019)
Knowledge distillation (Hinton et al. 2015)	Serving scale	Moves repeated inference cost into a one-time teacher-student training process	DistilBERT retains 97% of BERT accuracy with 40% fewer parameters and 60% faster inference (Sanh et al. 2019)

The table's three levers carry different deployment caveats. Pruning and quantization save energy only when the executed kernels skip arithmetic or move fewer bytes, which depends on sparsity

structure, kernel support, and conversion overhead (section 15.4.1.2). Knowledge distillation[20] instead amortizes its extra training cost across repeated serving.

Pruning, quantization, and distillation form the core toolkit for sustainable AI development, but their sustainability value depends on the measured bottleneck. The design question is not how to implement each compression method from first principles; it is how to rank these levers against memory movement, serving volume, carbon intensity, and lifecycle cost. A prioritized strategy orders them, together with carbon-aware scheduling, by the emissions term each one changes and attacks the largest term first.

Checkpoint 15.6: Prioritizing decarbonization strategy

Deploying a 70B LLM for a latency-sensitive application:

- ☐ **Technique ranking**: state the measurements needed to rank INT4 quantization, unstructured pruning, carbon-aware scheduling, and knowledge distillation for this service. Separate reduced energy demand from reduced carbon per unit of energy.
- ☐ **Systems justification**: propose a conditional ranking for a memory-bound service at fixed quality and latency, including hardware support, expected serving volume, extra training cost, and whether requests can move across regions or time.

15.9.2.2 Architectures that avoid unneeded work

Architecture is a lever beyond compressing a given model, because it sets how much work the hardware performs at all. Specialized datapaths and local data reuse can reduce energy per useful operation, but a comparison must match workload, precision, achieved quality, and measurement boundary. Chip arithmetic alone excludes memory and supporting systems; end-to-end energy also includes their activity and execution time.

Under a fixed power envelope, the cheapest joule is the one the hardware never spends, so the choice of which work an architecture performs at all bears directly on the ceiling. Alternative computing architectures expose potential sources of efficiency: selective activation, local state, and reduced communication. Comparing their energy with a dense model requires a shared task, quality threshold, latency target, and measurement boundary.

The architectural opportunity is to avoid work that the task does not need. A periodically evaluated dense model may repeat computation on unchanged inputs, while an event-driven design can respond only when a relevant signal changes. The comparison reduces to two rates. A dense model evaluated f_{eval} times per second at energy E_{eval} per evaluation draws $f_{\text{eval}}E_{\text{eval}}$ whether or not its input changed. An event-driven design that fires r_{event} times per second at energy E_{event} draws $r_{\text{event}}E_{\text{event}} + P_{\text{detect}}$, where P_{detect} is the constant power of the circuitry that watches for events. Skipping work pays only while $r_{\text{event}}E_{\text{event}} + P_{\text{detect}} < f_{\text{eval}}E_{\text{eval}}$, so the advantage shrinks as inputs change more often or as detection itself grows expensive. Dense execution can also reuse data and power down idle circuitry, so the comparison must examine the implemented workload rather than the architecture label alone.

Powering down inactive components helps either design only where the hardware supports it, and static power, sensing, and communication overhead remain part of the energy budget. Local learning and improved sample efficiency address different costs. Keeping an update local can reduce communication, while requiring fewer examples or optimization steps can reduce training work; neither property guarantees the other. Evaluation must include auxiliary preprocessing, adaptation, and inference costs at the required accuracy. Sparsity and sparse interconnects offer another route when execution skips arithmetic or transfers, but indexing, routing, load imbalance, and synchronization can consume the nominal saving.

Event-driven and sparse computation point toward research directions for sustainable AI. Architectures that implement **spiking neural networks (SNNs)**, event-driven models that communicate through discrete spikes, or sparse activation patterns can reduce energy. A neuron that does not spike neither computes nor sends a message, so switching and communication track input activity rather than a fixed clock[21] (Prakash et al. 2023).

[20] **Knowledge distillation**: Introduced by Hinton et al. (2015), distillation trains a compact "student" model on soft probability targets from a larger "teacher," capturing inter-class relationships that hard labels discard. The sustainability comparison must amortize teacher and student training across the actual serving volume and measure energy at the required output quality.

[21] **Spiking neural networks (SNNs)**: Sparse spiking activity can reduce communication and switching, but the complete comparison includes training, encoding, sensing, and the hardware used to execute the model. A saving measured on one temporal task does not transfer automatically to other workloads.

These directions expand the design space without making one architecture universally sustainable. The relevant improvement is less measured lifecycle impact per useful task at the required quality and service level. Dense, sparse, and event-driven implementations must meet the same comparison contract.

15.9.2.3 Lifecycle-aware systems

Repeated model updates add to the lifecycle footprint even when each deployed version is efficient. The relevant comparison is the work required to maintain the service's quality over time, including development and retraining. A retraining cycle that repeats the development sweep of section 15.5.2 pays its 5× final-run multiple again on every update. Transfer learning and fine-tuning can reuse a pretrained model instead of repeating training from scratch (Raffel et al. 2020), but the comparison must include adaptation, evaluation, and any additional work needed to recover quality on the target task. Reuse is a candidate saving whose value depends on that complete update cycle.

The deployment boundary matters just as much as the training boundary. Edge deployment can reduce communication energy by running inference on specialized low-power hardware at the point of use (Xu et al. 2020), but the lifecycle account must include the embodied carbon of manufacturing and replacing those devices. Embedding LCA methodologies into AI workflows allows teams to see this trade-off early: a cloud model, an edge model, and a hybrid cascade may have different winners depending on query volume, device lifetime, grid carbon intensity, and retraining frequency (International Organization for Standardization 2006a, 2006b). Henderson et al. (2020) argue for systematic reporting of ML energy and carbon footprints precisely so these comparisons can be made consistently. As the Jevons paradox warns, the accounting must also include usage growth, because optimizing one stage may increase total impact if lower costs enable wider deployment.

15.9.2.4 Benchmarks and operating metrics

Benchmarks matter when they make efficiency visible at procurement and design time. The ML.ENERGY Leaderboard (ML.ENERGY Initiative et al. 2023) reports inference energy measurements under declared model and serving configurations. MLPerf power submissions pair performance results with system-level power measurements under benchmark rules. These results support comparisons within the stated workload, quality, and measurement conditions; converting them into carbon footprints still requires electricity and lifecycle accounting beyond the benchmark.

The right efficiency metric depends on the serving regime. Batch inference is naturally expressed as correct samples per joule, latency-sensitive serving as queries per joule that meet the service-level objective (SLO), and generative workloads as joules per accepted output token at a declared quality threshold. Each is fleet energy productivity (equation 1.5) with useful work counted in the regime's own unit, and those queries are the requests that SLO goodput credits (section 10.3.1). Raw throughput per joule rewards systems that lower accuracy, truncate outputs, or miss latency objectives. Standardization does not make one platform universally green; it fixes useful work, quality, the measurement window, and the energy denominator well enough for procurement and architecture decisions to be argued quantitatively.

For sub-watt TinyML deployments, the MLPerf Tiny benchmark suite applies this discipline to constrained devices (Banbury et al. 2021). Tasks include visual wake words, keyword spotting, anomaly detection, and image classification. The closed division constrains implementations to reference-equivalent models, while the open division permits broader model choices. Reading accuracy, latency, and energy together within the declared division prevents a change in task or model quality from masquerading as a hardware-efficiency gain.

Energy and latency still have to be read together through the energy-delay product (EDP) (equation 15.16), which balances energy consumption against response time by penalizing solutions that save power only by taking too long:

$$\text{EDP} = E_{\text{total}} \times T = P_{\text{avg}} \times T^2 \tag{15.16}$$

where E_{total} is energy consumed, T is execution time, and P_{avg} is average power. The quadratic delay term penalizes solutions that achieve low energy through excessive delays. Lower EDP is preferable under this particular trade-off score, after quality and latency requirements have been satisfied. It does not authorize missing a service deadline or replace energy per completed task.

For TinyML deployments, EDP provides a score for comparing feasible operating points. Reducing clock frequency can lower power while extending execution; reducing voltage can lower switching energy as well, when the hardware permits it. Static power and the longer active interval can offset those gains. The preferred setting therefore comes from measured energy and delay across supported configurations, after excluding settings that miss the workload's deadlines.

Sustainability metrics complement performance benchmarks by recording the resource cost of useful work. Regulation is beginning to require such records. The EU AI Act's documentation obligations for covered general-purpose models, examined with the other regulatory mechanisms in section 15.10.1, include known or estimated energy consumption. Such requirements make the measurement boundary and provenance part of the reporting evidence rather than an optional annotation to a performance result.

15.9.3 Infrastructure optimization

Even the best-tuned training kernel cannot control how much power the chillers and transformers spend around the rack, or how much coal the regional grid burns to supply the substation. While algorithmic optimizations reduce operation counts, data movement, or the energy of execution, the surrounding operational environment determines the facility overhead and emissions associated with that work. Renewable energy integration, carbon-aware workload scheduling, and cooling optimization act on these surrounding terms, so each intervention must be evaluated at the boundary it changes.

A single hyperscale data center can consume over 100 MW of power, comparable to a small city.[22] Renewable energy integration, advanced cooling, and AI-driven optimization offer complementary interventions. Their value depends on whether the measured problem lies in the electricity source, heat-removal overhead, or operating controls.

Power availability and power source are connected infrastructure decisions, but annual energy matching does not establish continuous supply. Organizations training large models may locate near renewable sources or purchase renewable energy certificates (RECs) to claim the attributes of renewable generation. A certificate does not physically deliver electricity or remove the facility's demand from the grid; the design must still provide power when and where the workload runs.

Procurement can also support additional low-carbon supply. Microsoft's 2024 agreement with Constellation supported a planned nuclear-reactor restart, while Google's geothermal partnership with Fervo began supplying the Nevada grid in 2023. These examples address the need for generation beyond variable wind and solar output. They do not make one generation technology mandatory; the engineering requirement is a supply portfolio that meets the facility's demand across time, with transmission, storage, and workload flexibility accounted for.

Major cloud providers have announced renewable-energy commitments, but variable generation leaves a supply-matching problem. Storage and flexible scheduling can help address it alongside transmission and other generation sources; the appropriate combination depends on the site and workload. Google's published 2030 target for 24/7 carbon-free energy[23] illustrates the harder version of this goal: matching electricity consumption with carbon-free generation, including nuclear as well as renewables, hour by hour rather than only on an annual basis.

Cooling overhead depends on rack density, climate, and the heat-removal design, as quantified in section 15.6.3. Liquid cooling can support dense accelerator configurations, but its energy benefit depends on the complete pumping and heat-rejection system. Control policy supplies another lever. The DeepMind case in section 15.4.1.4 shows how measured operating data can guide cooling adjustments within thermal constraints; its reported savings describe the evaluated site rather than a universal reduction for any facility.

15.9.4 Carbon-aware placement and scheduling

A batch training job that consumes the exact same megawatt-hours can emit several times more carbon if it runs during an evening fossil-gas peak or in a coal-heavy region than if it waits for a windy afternoon or runs on a hydro-heavy grid. Placement and scheduling change the emissions attributed to a fixed workload rather than the energy it uses: placement chooses the region a job runs in, and scheduling chooses its hours. Both depend on the grid signal they act on and on the additional energy the move itself causes.

22 **PUE gap**: The PUE model in section 15.4.1.3 shows that facility overhead scales with every watt of IT load. At 100 MW, even a small PUE change therefore represents megawatts of cooling and power-delivery capacity, but its carbon value still depends on the grid signal.

23 **24/7 carbon-free energy (CFE)**: Hourly matching strengthens the time resolution of electricity procurement and accounting. Its physical impact depends on deliverability, additional clean supply, storage, and grid conditions. It does not eliminate manufacturing emissions or prove that every consumed electron came from a carbon-free generator.

↩ Scheduling against an external price signal, trading waiting and preemption against deadlines, is established in section 8.8.5.

15.9.4.1 Regional grid intensity data sources

Every carbon-aware decision inherits the quality of the grid signal behind it, and the right signal depends on the decision. Grid carbon intensity varies with generation mix, which changes hourly with demand, renewable availability, and plant dispatch. The US Energy Information Administration publishes historical grid emissions factors by region, updated annually. These averages give a coarse baseline for inventories or early planning, but they cannot resolve the hourly effect of moving a flexible workload. Electricity Maps provides time-resolved, flow-traced average and lifecycle intensity estimates across covered grids, while WattTime provides marginal operating-emissions signals that estimate the generator response to added or removed load.

The two kinds of signal answer the two questions that the measurement contract (section 15.4) separated. Time-matched average factors support the attributional inventory of a completed run. Marginal factors support the consequential question of how shifting or eliminating a run would change system emissions, which is the question a scheduler asks.

15.9.4.2 Placement across regions and hours

Carbon intensity varies across geographic locations and time periods, creating optimization opportunities. Temporal scheduling can reduce emissions when deadline-tolerant workloads are shifted toward lower-carbon hours or regions, with the realized gain depending on workload flexibility, grid mix, and whether the scheduler uses average or marginal emissions (Patterson et al. 2022; Radovanovic et al. 2021). **Carbon-aware scheduling** automates that shift, moving flexible, nonurgent work toward times and regions with lower carbon intensity while respecting its deadline and the available capacity.

A real placement decision weighs more than the grid factor. Adding electricity prices, facility overhead, and an internal carbon price turns the comparison into a placement objective, priced as a modeled total cost of ownership (TCO) over three candidate regions (the US-average grid, Quebec, and Iowa). Notebook 15.2 isolated the grid factor by holding facility energy fixed. This objective must instead let facility overhead differ from region to region, so it holds the run's IT energy fixed and lets each region apply its own PUE. That choice separates the carbon charge from the cost of facility overhead.

↩ The total-cost-of-ownership model that prices ownership, maintenance, and electricity is established in section 12.9.6.1.

Example 15.1: Automated carbon-aware placement

Scenario: A team is planning a large training run requiring 10,000 MWh of IT energy, which occupies an 8,192-GPU frontier-scale cluster for about 112 days. The team is choosing among three candidate regions (US average, Quebec, and Iowa) with different facility overheads (PUE) and grid carbon intensities.

Diagnosis: With an internal carbon tax of $100/tonne, the placement model adds a carbon charge to each region's modeled cost of ownership and selects Quebec for a total cost of $25.3M. Hardware ownership and maintenance make up most of that total and are the same in every region, so the candidates differ only in facility overhead and carbon. Without a carbon price, Quebec's lower PUE saves about $72,000 against the US-average placement. With the tax, the US-average and Iowa placements cost $0.5M and $0.8M more than Quebec, and most of each gap is the carbon charge.

Systems lesson: Carbon-aware placement adds an explicit carbon price to the fleet scheduling objective. The selected region minimizes the modeled cost among these candidates under the stated assumptions. A deployment decision must also check capacity, data transfer, deadlines, and reliability.

The placement example compares regions using a fixed annual grid factor for each candidate. Temporal scheduling adds another decision by allowing a flexible job to wait for lower-carbon hours within one region. These are distinct levers, and their relative value depends on the available regions and the job's deadline; section 15.9.4.3 develops the temporal lever with the marginal signal that the measurement contract assigns to it.

15.9.4.3 Time-shifting under deadlines

Carbon-aware scheduling is a load-shifting software problem. The scheduler queries one of the grid signals compared above and combines it with workload flexibility. A feasible shift must finish before the deadline, retain access to data and accelerator capacity, and preserve the job's reliability constraints. Its net benefit must exceed checkpoint, transfer, restart, and idle-capacity emissions. Pausing useful work without reducing the cluster's powered load can erase the expected saving.

Google's carbon-intelligent computing platform illustrates this approach by using forecasts and virtual capacity limits, hourly caps on the flexible compute a cluster may run that are set lower in high-carbon hours, to delay flexible work (Radovanovic et al. 2021). In the multi-tier accounting cascade of figure 15.17, temporal load shifting modulates the grid emissions factor applied to facility electricity rather than reducing raw compute FLOPs. That is a carbon intervention, not automatically an energy reduction; checkpointing, transfer, and a different runtime can also change electricity consumption.

Because a scheduler asks how moving a job changes emissions, it needs a marginal signal rather than the average factor an inventory uses, and the providers compared above expose signals with different methodologies and coverage. A feasible schedule must satisfy three constraints on one time axis at once, the grid carbon signal $CO_{2,\text{grid}}(t)$, the job's execution state, and cluster headroom $N(t)$ (figure 15.19). The job waits, checkpointed, through the high-carbon peak, runs for T_{job} in a low-carbon trough where headroom exists, and still finishes with slack before its deadline. Waiting pays only when the avoided grid emissions exceed the checkpoint I/O, idle standby, and state migration that the wait itself causes.

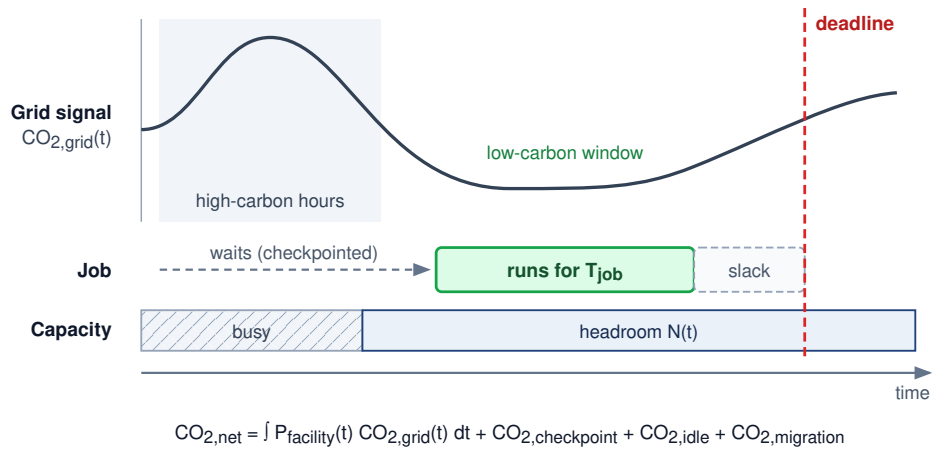

Figure 15.19: A Feasible Carbon-Aware Schedule: Three constraints share one time axis. The job waits, checkpointed, through the high-carbon hours of the grid signal and runs in the low-carbon window only where cluster headroom exists, finishing with slack before the deadline. The equation below the timeline charges the shift for its checkpoint, idle-wait, and migration carbon, so waiting pays only when the avoided emissions exceed those terms. The curve is schematic, not a measured grid trace.

A large distributed training run draws sustained aggregate power with the correlated phase transitions described in section 15.6.2, but its completion deadline may leave scheduling flexibility. Checkpoints created for fault recovery also provide the restart points for the deliberate pause in figure 15.19. The decision requires a local emissions forecast and enough deadline slack; late-afternoon or overnight operation is not universally preferable. Fault tolerance thus supplies a mechanism for flexibility without proving that every pause saves emissions. Online inference has tighter delay budgets and may support bounded preemption, while offline inference can often be deferred more freely.

The **duck curve** describes the steep evening ramp in net load (total demand minus variable solar and wind output) as midday solar generation falls toward sunset while residential and commercial

↩ Checkpoint-restart mechanics and the failure-driven choice of checkpoint interval are established in section 7.8.

demand rises. Because grid operators typically meet that rapid net-load climb by dispatching fast-ramping fossil-gas generators, the evening ramp often coincides with a sharp spike in marginal carbon intensity. A constant data-center load raises the baseline level of net demand but does not by itself increase the evening ramp rate. A flexible training load, by contrast, can reduce both the required ramp and its own emissions by curtailing during that interval. Schedulers should still act on local forecasts and marginal emissions signals rather than a fixed daily pause window, which would ignore regional generation and transmission conditions.

Renewable generation variability complicates batch training dispatch across power grids. Over one week of November 2023, the European Union's hourly generation mix moved across days as well as within each day (figure 15.20). Demand peaked each day between 326 GW and 367 GW. Nuclear held near 73 GW, solar contributed midday output, and wind ranged from 27 GW on Sunday to 111 GW on Thursday. Wind contributes to that variation, but demand, other generation, storage, and electricity exchanges also affect fossil dispatch. Carbon-aware scheduling therefore needs regional grid forecasts rather than a fixed assumption about which hours are cleanest.

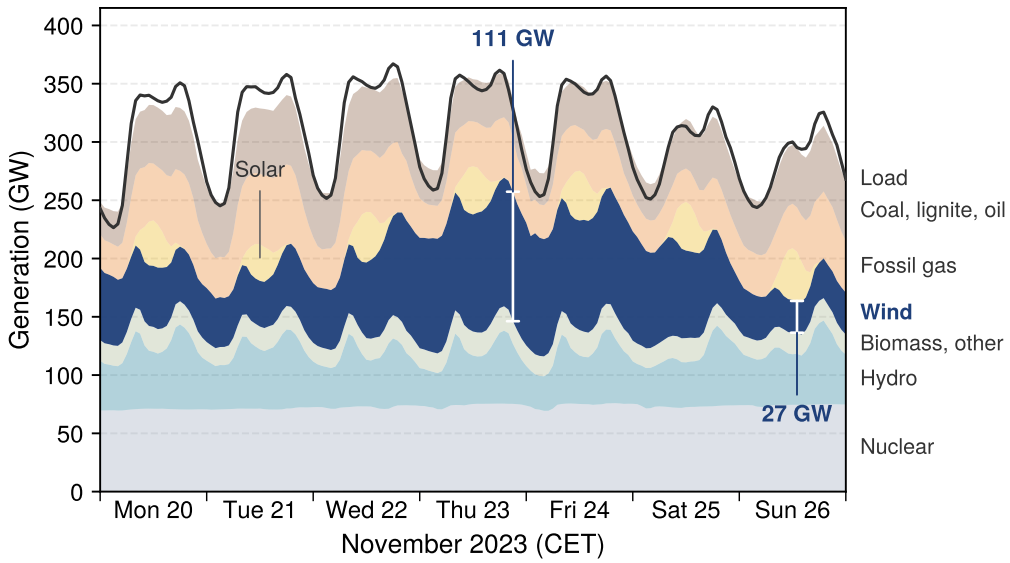

Figure 15.20: European Electricity Generation over One Week: Hourly public net generation by source for the European Union in week 47 of 2023 (20 to 26 November), stacked, with the load curve on top. Nuclear forms a nearly flat base, hydro follows the daily demand cycle, and solar adds only small midday bulges this late in the year. Wind swings across days rather than hours, from about 27 GW on Sunday to about 111 GW on Thursday, so the fossil gas and coal that fill the remaining gap move from day to day. Lower-carbon windows for deferrable training therefore come from the forecast for a specific week and region, not from a fixed daily pattern. Source: Energy-Charts (Fraunhofer ISE) public power data, from ENTSO-E and national statistics.

15.9.5 Google sustainability framework case study

When Google engineers compared a baseline Transformer training setup against one that combined a sparser architecture, custom TPUs, an efficient cloud facility, and a low-carbon grid, they reported an 83-fold reduction in energy and a 747-fold reduction in operational CO_2 emissions (Patterson et al. 2022). The value of Google's case study is that it groups mitigation into four levers that correspond to terms this chapter has tracked, namely the model's work, the machine's efficiency, the facility's overhead, and the grid's carbon factor. Google engineers identified these "4 Ms" (table 15.13) for reducing the carbon footprint of rapidly expanding AI workloads (Patterson et al. 2022).

Table 15.13: Google's 4 Ms Sustainability Framework: The framework combines model architecture, hardware specialization, infrastructure efficiency, and carbon-aware placement, each acting on a different term of the footprint. Their combined effect requires a consistent end-to-end account: component factors can be multiplied only when they describe compatible workloads and measurement boundaries.

Lever	Intervention	Reported efficiency effect	Example mechanism
Model	Select efficient AI architectures such as sparse models or neural-architecture-search-derived designs	5–10× lower computation requirements without compromising model quality	Evolved Transformer and Primer
Machine	Use AI-specific hardware rather than general-purpose systems	2–5× performance-per-watt improvement; TPUs show 5–13× greater carbon efficiency relative to nonoptimized GPUs	Tensor Processing Units
Mechanization	Run optimized cloud infrastructure at high utilization	1.4–2× energy reduction compared to conventional on-premises data centers	Lower facility PUE than the industry-average baseline used in the 2021 study
Map	Place data centers in regions with low-carbon electricity supplies	5–10× lower gross emissions	Real-time monitoring of renewable energy usage across infrastructure

In Patterson et al. (2022), that 83-fold energy reduction and 747-fold CO_2 reduction come from a consistent accounting model supported by measurements across compatible workloads and boundaries; multiplying unrelated best-case component ratios would not establish the same result. In the period studied, AI workloads also represented less than 15 percent of Google's total energy consumption despite expanding deployment. That share is evidence about relative consumption. It does not by itself establish constant or declining absolute energy when the denominator also changes, so engineers must assess absolute growth from measured energy totals alongside the reported efficiency improvements.

Empirical case studies demonstrate how sustainable-AI engineering can improve both capability and environmental impact. For example, Patterson et al. (2022) compare GPT-3 with Google's GLaM and report improved quality metrics alongside reduced training computation and lower-carbon energy sources, while the GLaM model paper explains the mixture-of-experts architecture behind that efficient scaling (Du et al. 2022).

The strategy in the case study, which combines systematic measurement, carbon-aware development, transparency in reporting, and renewable energy transition, provides a practical template for sustainable AI scaling. The analysis also shows that earlier extrapolations overstated ML energy requirements by large factors when they ignored empirical efficiency gains and measurement boundaries. Technical levers alone, however, only save carbon when organizations have an incentive to deploy them and a bound that prevents rebound growth, which brings the engineering account to external policy, regulation, and governance.

15.10 Policy, Regulation, and the Path Forward

Electricity prices reward some efficiency gains, but they need not capture water stress, manufacturing emissions, or local environmental burdens. A training cluster can reduce its electricity bill by moving to a cheaper region while increasing its emissions. If the bill omits that environmental cost, a cost-minimizing placement decision may favor the higher-emissions site. Reporting requirements expose the trade-off; carbon pricing, procurement conditions, and other policy instruments can change the incentives. Their effect depends on which activities they cover and how organizations implement them.

For a systems reader, policy is the control plane that changes the objective function. Reporting rules make hidden energy and embodied-carbon costs measurable, carbon pricing turns location into a scheduling variable, and procurement standards make lifecycle accounting part of infrastructure design. The regulatory mechanisms in section 15.10.1 matter because they determine which optimizations become economically rational at fleet scale.

15.10.1 Regulatory mechanisms

An organization training a frontier model above 10^{25} FLOPs or operating a large European subsidiary can face legal requirements to disclose energy consumption and value-chain emissions, yet neither law requires metering every individual AI job. Reporting rules act at different organizational boundaries: the EU AI Act[24] includes model-level technical-documentation obligations, while the **Corporate Sustainability Reporting Directive (CSRD)**[25] addresses sustainability reporting by covered companies. Model energy records can inform both, so engineers must identify the covered entity or model, applicable dates and exemptions, and the evidence the reporting method requires.

Policy can also impose a priced resource constraint. A hypothetical emissions cap for AI operations would specify tonnes of CO_2e and allow trading of unused emissions allowances, following the general logic of **emissions trading systems (ETS)**.[26] A power cap in MW or an energy budget in MWh would constrain different quantities and must not be called the same credit. The EU **Carbon Border Adjustment Mechanism**, by contrast, applies to embedded emissions in specified imported goods and electricity; it does not generally price an AI service by its compute location.

Incentives and procurement conditions can change which designs organizations choose. A grant can support efficiency research, while a procurement contract can require declared energy measurements or lifecycle evidence. These conditions need measurable criteria rather than an undefined "green" label. Scope 3 emissions also depend on whose inventory is being prepared: purchased hardware and outsourced cloud services enter the customer's value-chain account, while the supplier classifies its own activities within its organizational boundary. The same physical emissions must not be added twice merely because two organizations report them.

15.10.2 Industry self-regulation and standards

A cloud provider can announce "100 percent annual renewable matching" while its data centers still draw fossil-generated electricity every night when solar output drops to zero. Voluntary commitments and common standards supplement government requirements, but an announced target is not evidence of a measured reduction. Annual renewable-energy matching and 24/7 carbon-free energy (CFE) address different time resolutions: hourly matching makes periods of unmatched demand visible, while the resulting physical emissions still depend on the generation and grid context (Monyei and Jenkins 2018). Evaluating a provider's progress therefore requires the reported method and results, rather than merely its announced target.

Internal carbon pricing gives emissions a weight in project decisions. A "shadow price" lets an organization compare alternatives using both financial cost and attributed emissions; its influence depends on the price, accounting boundary, and whether decision makers use it. Measurement tools supply inputs to that comparison. Projects like **CodeCarbon** and **ML CO_2 Impact** let developers estimate model carbon footprints within their workflows (Schmidt et al. 2021; Lacoste et al. 2019), but the estimates still need declared power, location, and allocation assumptions.

Standardized benchmarks provide evidence for particular efficiency claims. **MLCommons** defines MLPerf power-measurement protocols that connect system energy with workload performance under declared conditions. The Green Software Foundation's Software Carbon Intensity specification addresses a different boundary, combining operational and allocated embodied emissions per functional unit. A benchmark result or a calculation method is not itself an independent sustainability audit. Accountability requires checking the reported inputs, scope, and comparison against the claim being made.

15.10.3 Public engagement and transparency

A sustainability report that trumpets a 50 percent drop in energy per query while omitting a fourfold expansion in total cluster capacity, or that leaves hardware manufacturing out of its inventory, masks a growing footprint behind a selective boundary. Misleading presentation of environmental performance, or **greenwashing**,[27] can arise from selective boundaries or omitted trade-offs as well as unsupported offset claims. Effective AI sustainability governance requires evidence that matches the claim, combining workload telemetry with facility measurements and supply-chain inventories. Runtime power counters alone cannot establish Scope 1–3 completeness or the environmental costs of manufacturing.

24 | **EU AI Act (2024)**: The Act presumes high-impact capabilities above 10^{25} FLOPs of training compute and imposes additional obligations on general-purpose models with systemic risk. Annex XI includes computational resources and known or estimated energy consumption in required technical documentation; qualifying open-source models without systemic risk can be exempt from these documentation obligations. Article 101 permits fines for specified provider violations up to 3 percent of annual worldwide turnover or EUR 15 million, whichever is higher. Application dates, model scope, and exemptions must be checked before turning these provisions into an engineering reporting requirement.

25 | **Corporate Sustainability Reporting Directive (CSRD)**: Directive (EU) 2026/470 narrows reporting thresholds to more than 1,000 employees and more than EUR 450 million in net annual turnover, assessed at group level where applicable. Member States must transpose its reporting amendments by March 19, 2027; entity scope, exemptions, and reporting-year rules still determine applicability. Covered reporting follows European Sustainability Reporting Standards rather than a separate inventory of AI jobs. The engineering contribution is traceable information on relevant electricity, procurement, and value-chain impacts.

26 | **Emissions trading for compute**: Applying cap-and-trade to AI emissions is a policy analogy here, not a description of an established compute-allowance market. Such a design would need a declared emissions boundary, measurement method, enforcement mechanism, and rules for trading allowances.

Pledge-style disclosure is useful only when it creates auditable data rather than reputational cover. The **Montréal Carbon Pledge**, originally a commitment by institutional investors to measure and disclose carbon footprints annually, is a useful model precisely because its value lies in the disclosed data, not the pledge itself. The same standard applies to an AI organization, whose sustainability claims are only as credible as the workload-level measurements behind them.

15.10.4 Environmental justice and responsible development

When a hyperscale data center is sited in a drought-prone agricultural basin or beside a neighborhood served by peaking gas turbines, local residents absorb the water stress and air pollution while the model's outputs serve users thousands of miles away. Transparency establishes the evidence base for those impacts; environmental justice asks how the burdens and benefits revealed by that evidence are distributed. Site selection is where much of that distribution is decided, which makes it an environmental-justice decision rather than a facilities decision alone. Environmental sustainability extends ML systems responsibility from model behavior to ecological stewardship (Vinuesa et al. 2020). The computational resources required for AI development concentrate environmental costs on specific communities while distributing benefits unequally across global populations. Data centers impose both grid demand and cooling-water demand, and the local burden depends on the energy mix, watershed, and competing regional uses.

That geographic concentration creates environmental-justice risks that align with broader responsible AI frameworks.[28] The fairness claim here is narrow and systems-specific. An ML fleet allocates benefits, electricity demand, water use, land pressure, and e-waste across different communities. Communities hosting AI infrastructure can bear disproportionate environmental burdens while having limited access to AI's economic benefits.

These distributional concerns fall into four groups (figure 15.21). Environmental justice concerns local water withdrawal and grid strain. Intergenerational equity concerns atmospheric CO_2 persistence and long-lived infrastructure. Resource competition concerns agricultural and municipal claims on land, water, and minerals. The transparency gap concerns undisclosed per-model energy, carbon, and water records. Engineering acts most directly on the last group, because auditable telemetry is what allows the other three to be assessed at all; the measurement contract of section 15.4 supplies that telemetry.

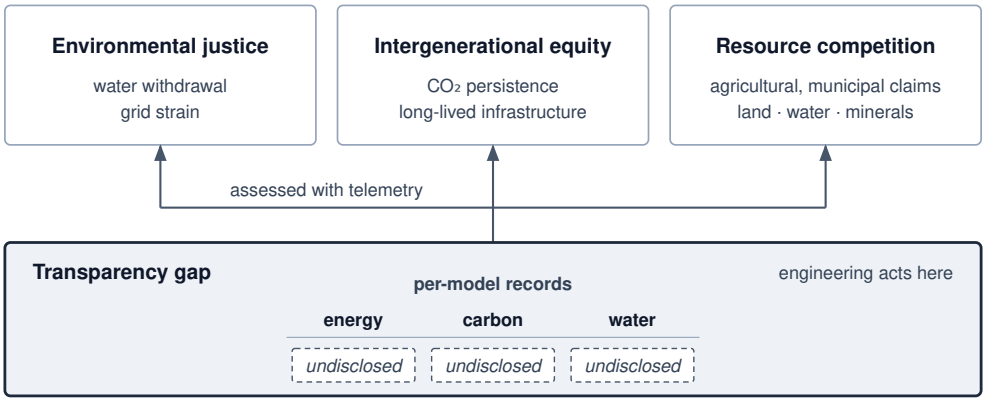

Figure 15.21: Ethical AI Concerns: The distributional concerns raised by AI infrastructure fall into four groups. Environmental justice, intergenerational equity, and resource competition stand on the transparency layer, because each can be assessed only with the per-model energy, carbon, and water records that the transparency gap leaves undisclosed. That layer, drawn with the heavier outline, is the one engineering telemetry addresses most directly.

Impact assessments and community participation make these trade-offs visible before a site is committed. Any claim that such an assessment is legally required must identify the applicable jurisdiction, rule, and project scope. Where siting and governance bound the burden of current facilities, longer-term research targets the architectural and measurement limits that make AI workloads so resource-intensive in the first place.

27 **Greenwashing in AI**: Typical forms are claiming "carbon neutrality" through offsets while expanding data center capacity, or highlighting per-query efficiency gains while total compute grows faster. Green-claims rules can require verifiable evidence for environmental claims. For ML engineers, the technical litmus test is whether sustainability reporting covers all three GHG Protocol scopes, or conveniently omits Scope 3 (hardware manufacturing, cloud supply chain) where much of AI's carbon can reside.

28 **Environmental justice in data center siting**: A site-cost model can omit burdens borne by nearby residents, such as competing water demand or changes to local infrastructure. A low financial cost therefore does not establish a low community burden. Site evaluation must identify who benefits, who bears the resource costs, and how affected communities participate in the decision alongside the analysis of grid intensity and water stress.

15.10.5 Future research directions

Even after a team applies every scheduling, cooling, and compression lever in this chapter, conventional accelerators still spend most of their dynamic energy shuttling weights and activations back and forth across memory buses, while most teams still lack real-time, per-job carbon telemetry. The research agenda therefore follows the same engineering logic as the mitigation framework: reduce data movement, close the measurement gap, and avoid redundant computation. **Non-von Neumann computing architectures**,[29] including **neuromorphic computing** and **in-memory computing**, can reduce selected transfers, extending the event-driven and spiking designs of section 15.9.2.2 from avoided work to avoided movement. Evaluate the remaining arithmetic, conversion, control, and peripheral energy at the required quality; eliminating one cost does not eliminate the whole system account.

[29] **Von Neumann bottleneck**: Separating processing from memory requires data transfers whose energy depends on locality and arithmetic intensity. If movement contributes a fraction q of total energy, eliminating that term alone gives at most a $1/(1-q)$ reduction factor when the remaining terms are unchanged. For an illustrative 60–80 percent share, that bound is 2.5–5×, before accounting for replacement circuitry.

A critical implementation barrier is the "measurement gap": AI teams need standardized, workload-level reporting of energy use and carbon emissions rather than relying only on coarse proxy metrics (Henderson et al. 2020). Coarse methods often rely on proxies such as GPU-hours multiplied by average grid intensity, which fail to capture the real-world dynamics required by reporting regimes that care about time, location, and workload boundary. Developing and standardizing granular, real-time energy and carbon accounting tools is essential for both compliance and effective optimization.

A second research direction reduces redundant computation before it reaches the accelerator. Data selection, model design, and training procedure all affect the environmental cost of learning (Wu et al. 2022). Smarter data sampling, active learning, and data valuation can reduce work spent on low-value examples, but the resulting model quality and coverage must be measured rather than assumed. An integrated approach combining algorithmic efficiency, hardware innovation, renewable energy adoption, and transparent governance is necessary to align AI development with environmental constraints.

Reducing redundant computation can improve operational efficiency, but compliance and environmental benefit require their own evidence. A sustainable design must show that the saved work survives the complete lifecycle account and the resulting deployment scale. Incorrect boundaries or assumptions can otherwise turn a local improvement into a larger total footprint.

15.11 Fallacies and Pitfalls

Sustainability involves counterintuitive physics where efficiency improvements can increase total consumption and geographic choices can outweigh component-level optimizations. These fallacies and pitfalls capture errors that waste compute budgets and planetary resources through misallocated optimization effort.

Fallacy: *Cloud computing automatically makes AI systems more environmentally sustainable.*

Engineers assume cloud providers operate efficiently and sustainably. Cloud infrastructure can improve utilization and facility efficiency, but it transfers neither the physical footprint nor the customer's accounting responsibility out of existence. The 7B model trained on 64 A100s in notebook 15.5 emits roughly 21.4× as much CO_2 on the US-average grid as on Quebec's hydroelectric grid, and the hydroelectric figure is an operational factor, not a lifecycle-zero one. As demonstrated in section 15.5, cloud comparisons must disclose region, energy boundary, utilization, embodied allocation, and whether the electricity is Scope 2 for the provider or Scope 3 for the customer.

Pitfall: *Focusing only on operational energy consumption while ignoring embodied carbon and lifecycle impacts.*

Teams optimize training efficiency while ignoring manufacturing emissions. In low-carbon grids, embodied carbon can dominate fleet-level footprint accounting. As quantified in section 15.5.5, an A100 accelerator carries roughly 150 kg CO_2 from manufacturing (Luccioni et al. 2023); for the earlier 14-day, 64-A100 training run, the unamortized upfront burden is roughly 9.6 t CO_2. Amortized over a 4-year service life, the share attributable to this specific job is about 90 kg. The job comparison must use that allocated share; a fleet lifecycle comparison must instead compare the full manufacturing burden against operations over the same fleet lifetime. Comparing the entire manufacturing burden with one short job's electricity would mix boundaries. Extending hardware lifetime from 3 years to

5 years reduces amortized embodied carbon by 40 percent, but procurement decisions must also evaluate the operational cost of retaining the hardware.

Fallacy: *TDP is actual power consumption.*

Thermal design power is a hardware design specification, not the measured power of a particular workload. Actual draw depends on the device, configured power limit, utilization, memory traffic, clocks, and idle intervals. A TDP-based calculation is a stated capacity scenario; it does not establish serving energy or facility energy. Carbon and electricity-cost estimates that drive placement decisions (section 15.5) should use measured power over the workload interval and include the support systems within the chosen boundary.

Pitfall: *Using PUE as a complete environmental metric.*

Data-center operators report PUE in sustainability disclosures and engineers may treat it as a complete environmental score. PUE measures only facility energy relative to IT energy. A lower-PUE facility can still cause more electricity-related emissions if its grid factor is sufficiently higher, and neither factor accounts for hardware manufacturing or local water stress. Compare matched emissions factors and complete lifecycle boundaries rather than ranking sites by PUE alone.

Fallacy: *Efficiency improvements automatically reduce total environmental impact.*

Engineers assume that halving inference cost cuts environmental impact in half. The Jevons paradox warns that efficiency improvements can increase total consumption by enabling expanded usage. In a rebound scenario, suppose reducing token cost from $0.06 to $0.002 per 1,000 tokens accompanies a proportional 30× reduction in measured energy per query at fixed output length and quality. A 100× increase in query volume then raises total energy by 3.3× and, at unchanged grid intensity, emissions by the same factor. Any implementation change that lowers per-query energy can still raise total energy when usage grows by a larger factor. Teams must therefore evaluate deployment volume against the intended carbon or resource budget. Per-query gains alone do not control the rebound illustrated in figure 15.18, where the red area of new consumption at the lower unit cost completely swallows the green area of initial efficiency savings.

Pitfall: *Treating carbon offsets as a substitute for reducing actual emissions.*

Organizations may purchase offsets without establishing that the credited reduction is additional, durable, independently verified, and free of double counting. An offset does not change the electricity consumed by the workload or its immediate emissions. Teams should report physical emissions reductions separately from compensation claims, with the offset boundary and uncertainty visible. Geographic placement can reduce electricity-related emissions when the complete workload comparison supports it; neither an offset purchase nor a renewable label substitutes for that calculation.

Fallacy: *Component-level optimization guarantees lifecycle improvement.*

Teams reduce training cost to improve sustainability without analyzing deployment scale. A training change that increases serving energy must be evaluated over expected query volume: the crossover occurs when accumulated additional serving emissions exceed the one-time training saving. Edge deployment adds device manufacturing and replacement to any avoided data-center and transmission emissions. Keeping hardware longer avoids new manufacturing but may consume more operating energy than replacement. Each comparison requires the same useful service, grid-factor boundary, and lifetime; there is no universal query count or upgrade age that settles it. Effective sustainability requires the complete account in section 15.5.6.

Pitfall: *Evaluating sustainability at the component boundary instead of the lifecycle boundary.*

A pruned model may contain fewer weights while its sparse representation adds indexing and execution overhead. If those costs outweigh the arithmetic and transfers avoided, the smaller model need not consume less energy. The same boundary error appears when hardware teams report accelerator efficiency without procurement carbon, platform teams report PUE without grid intensity, or model teams report training emissions without expected serving volume. The lifecycle boundary in section 15.5.6 requires evaluating training, serving, embodied carbon, regional grid mix, hardware lifetime, and demand growth together before calling a change sustainable.

15.12 Summary

The lifecycle boundary connects training, serving, embodied carbon, grid mix, hardware lifetime, and demand growth. Sustainable AI asks whether the machine learning fleet can operate within available energy, water, and material resources even after its models and infrastructure have been optimized.

The measurement contract makes comparisons meaningful by fixing useful work, required quality, system boundaries, attribution, time, location, and uncertainty. Operational and embodied carbon accounting then identify which lifecycle term an intervention can change. Memory traffic, idle equipment, cooling, and power delivery can all consume energy beyond useful arithmetic; lower FLOP counts alone do not establish lower environmental impact.

The Fleet Law connects execution time to the fleet energy account. A step's energy is approximately its average fleet power times $T_{\text{step}}(N)$ (equation 15.13), so exposed communication and synchronization extend the interval over which the system draws power. Recovery adds further work outside successful steps. The C^3 terms therefore matter to both performance and sustainability, although the power trace determines their energy cost. PUE extends an IT-energy estimate to the facility boundary, the grid factor attributes operational emissions, and embodied carbon adds the hardware lifecycle burden rather than converting operating joules.

Mitigation follows that diagnosis. Algorithm and hardware changes must reduce measured resource use while preserving quality, latency, and reliability; scheduling must account for grid intensity and the cost of moving work; hardware replacement must weigh operational savings against new manufacturing emissions. Water availability and the distribution of environmental burdens also matter when choosing where and how a fleet can grow.

Key Takeaways: Efficiency alone is not enough

- **Power is a hard ceiling**: A fleet cannot compute past the megawatts, cooling, water, and grid capacity its site can supply. Sustainability is therefore an existence constraint on model scale, not a public-relations layer around an otherwise finished architecture.
- **A footprint is only as good as its boundary**: An energy or carbon figure supports a comparison only when it names the functional unit, system boundary, attribution rule, carbon signal, and time and location. The boundary alone can more than double the energy reported for the same prompt.
- **Demand can outrun efficiency**: Lower energy per query does not ensure lower fleet energy use. At fixed work per query, total energy rises when query volume grows by a larger factor than per-query energy falls, so efficiency targets need an absolute resource budget.
- **Memory traffic can dominate serving energy**: During bandwidth-bound autoregressive decode, accelerators can draw power while waiting for weights and state. Quantization, batching, or memory-optimized hardware help sustainability only when they reduce measured energy at the required quality, latency, and reliability.
- **Carbon starts before boot**: Hardware manufacturing creates emissions before the first query runs, and its lifecycle share rises on clean grids and under low utilization. Keeping or replacing hardware therefore requires comparing new manufacturing emissions with operational savings over the same future service.
- **Location changes the footprint**: Moving flexible work to a cleaner grid can reduce operational carbon only when the saving exceeds transfer, restart, and idle-capacity emissions. Placement must also preserve deadlines, data access, capacity, and reliability; regional intensity ratios alone do not establish the net benefit.
- **Water is its own account**: Evaporative cooling saves facility energy by consuming local water, while dry cooling saves that water by spending electricity whose generation consumes water elsewhere. Report withdrawal and consumption, direct and indirect, weighted by the stress of each watershed.

Sustainability changes the scale at which efficiency must be judged. A reduction per query is a local result; the fleet result depends on how many queries run, where they run, and what hardware their service requires over time. Two forces push demand past efficiency. Scaling-driven growth is the first, as scaling laws (principle 9) price each further quality gain in more compute per model. The Jevons paradox of AI (principle 25) is the second, as cheaper work invites more of it. Both make the distinction between efficiency per task and absolute demand consequential, because efficiency creates room within a resource budget while demand determines how much of that room is consumed. Deployment decisions therefore need both a useful-work efficiency target and an absolute resource budget.

What's Next: From sustainability to responsibility

A fleet that is secure, robust, and within its power budget is not automatically fit to deploy. Security, robustness, and sustainability each answer a technical question about what a production system can support. Section 15.10.4 already showed that a feasible fleet can place its water, grid, and waste burdens on communities that see little of its benefit. The model's decisions raise the same question about the people they act on. An accurate, fast model can still concentrate its errors among people with little power to challenge them. Responsible AI therefore asks who receives the system's benefits, who bears its costs, and how an affected person can contest a decision. The carbon inventory and lifecycle account built here become one input to the assurance case that answers it. They supply evidence for that judgment; they do not settle whose interests the fleet should serve.

↪ Governance frameworks, algorithmic fairness, and ethical guardrails are developed in chapter 16.

16

Responsible AI

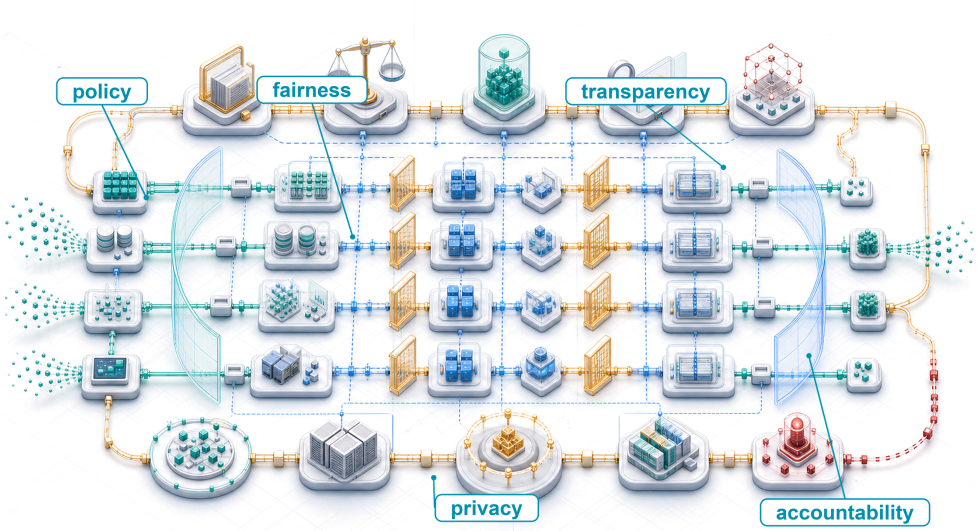

16.1 The Governance Imperative
16.2 The Governance and Assurance Lifecycle
16.3 Core Principles and the ML Lifecycle
16.4 Architectural Trade-Offs
16.5 Fairness and Bias Mitigation
16.6 Privacy and Data Governance
16.7 Explainability and Transparency
16.8 AI Safety, Robustness, and Generative Controls
16.9 Accountability, Governance, and Sociotechnical Dynamics
16.10 Governing Responsible AI at Fleet Scale
16.11 Fallacies and Pitfalls
16.12 Summary

Purpose

Why can an accurate, fast model still be unfit for deployment?

In high-risk regulated settings, a model that cannot explain or document its behavior, support audit, or meet applicable discrimination rules may fail deployment review regardless of its accuracy or latency. Strong aggregate performance can conceal a pattern of errors concentrated among people with little power to challenge the outcome. Increasing throughput then increases the reach of that failure. A single-node evaluation can measure predictions, but it cannot establish whether every deployed version uses an approved decision policy or whether an affected person can obtain a correction. Across a fleet, records, models, interfaces, and organizational authority must remain connected as each changes. An explanation delivered too late may be useless to someone facing an immediate decision; retaining every detail indefinitely can create a separate privacy problem. Investigating an alert requires staff and evidence, while correcting a decision may require action across several dependent services. These obligations consume resources and can constrain deployment even when inference capacity is ample. They also persist after launch: a change in population, workflow, or model behavior can invalidate the evidence that originally justified operation. Responsible AI engineering makes those obligations part of the organization's daily operation, with resources and authority to act when measured behavior falls short. In C^3 terms, responsibility is a coordination constraint: the fleet needs evidence and authority to decide when compute may run, when it must stop, and who can correct its consequences.

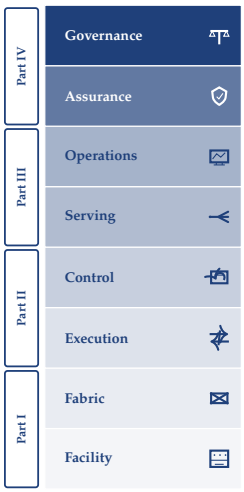

Learning Objectives

- Explain why an accurate, fast model can still fail deployment review, and why each safeguard is budgeted capacity whose placement decides who is protected
- Calculate fairness metrics and evaluate observational versus causal definitions under base-rate and prediction trade-offs
- Design bias detection and fairness monitoring across demographic groups, drift, and fleet-scale feedback loops
- Evaluate privacy and machine-unlearning controls against accuracy, compute, and governance requirements
- Select explanation methods under latency, compute, auditability, and user-contestability constraints
- Analyze human-AI feedback loops, automation bias, and value conflicts in deployment contexts
- Construct an auditable assurance case, including bounded robustness and generative-system controls, from intended use and evidence through launch, remedy, material change, and retirement

16.1 The Governance Imperative

In 2018, Reuters reported that Amazon had abandoned an experimental hiring algorithm trained on historical resume data after discovering it systematically penalized female candidates (Dastin 2018). The historical records supplied a computable training objective, but reproducing patterns in those records was not the same as identifying qualified applicants fairly. The experiment exposed a mismatch between the learned objective and the legitimate hiring goal, because historical success could encode unequal access rather than merit-based qualifications.

That mismatch between a working pipeline and a legitimate outcome sits beyond the traditional systems agenda of architecture, infrastructure, data processing, security, and robustness (Stoica et al. 2017). The preceding chapters made the fleet secure against adversaries, robust under shift, and accountable for the energy it draws. The siting analysis in section 15.10.4 left one question that energy and carbon accounting could expose but not settle, namely who bears the environmental cost of where the fleet runs. Responsible AI carries that question from the fleet's resource footprint to the decisions the fleet makes about people. As the hiring failure shows, a service that remains available, secure, accurate, and within its resource budget can still distribute harms unfairly or deny affected people a remedy. Fairness, transparency, and accountability therefore need measurable requirements and operating controls alongside the mechanisms that keep the fleet running.

Across the multi-tier fleet hierarchy in figure 1.4, responsible AI completes the Governance band, the point where the system meets the real world. The Infrastructure, Distribution, and Serving bands beneath it supply compute, data movement, distributed execution, and serving infrastructure. Within the Governance band, the Assurance row contributes security controls and robustness mechanisms, and the Governance & Sustainability row contributes sustainability budgets. Responsible AI weighs that evidence and adds the constraints that determine why the machine runs and whom it serves. It requires owners to identify, assess, limit, and remedy harms from an otherwise correctly functioning system. Alongside the Fleet Law's accounting of execution time, responsible AI adds conditions for continued operation: evidence that the deployment remains within its approved bounds and authority to intervene when it does not. Toxic outputs can erode user trust, discriminatory decisions can distort future training data when those decisions shape data collection, and privacy violations can require restricting or suspending a service. These consequences depend on the deployment, but the obligation to investigate harm does not depend on whether it also degrades model performance. Responsible AI therefore shapes objective selection and operating constraints from the beginning of design.

↩ The Fleet Law governing execution time and scaling efficiency is established in section 1.5.4.

Definition 16.1: Responsible AI

Responsible AI is the practice of designing, auditing, and operating ML systems to measurable fairness, safety, privacy, and accountability standards, translating ethical principles into verifiable system properties that constrain model training, deployment decisions, and operational monitoring.

1. **Significance**: Responsible AI constraints impose real costs: fairness-aware training can lengthen training runs, real-time bias monitoring can consume serving-path latency, and on-demand explainability can require additional model evaluations or asynchronous worker capacity. At fleet scale, these safeguards must be budgeted like reliability or security controls. Conversely, a facial recognition system with a persistent subgroup error gap can affect millions of users before detection, creating regulatory and liability costs that dwarf the monitoring investment.
2. **Distinction**: AI ethics examines normative principles about what systems should do; responsible AI engineering turns selected requirements into controls and evidence. Bias detection algorithms, differential privacy implementations, audit trails, and architectural guardrails support bounded claims that reviewers can inspect. Choosing legitimate requirements and judging the remaining harms still require domain expertise, stakeholder participation, and accountable authority.
3. **Common pitfall**: A frequent misconception is that responsible AI is a final compliance review applied to a finished model. Postprocessing, which adjusts a trained model's decision thresholds, can satisfy selected metrics on observed outcomes, but threshold adjustment cannot repair an invalid target, unreliable labels, or missing information. Those defects may require changes to data collection, labels, or learned representations. The remedy depends on the source of harm; a final threshold check cannot substitute for investigating it.

That practice must account for **sociotechnical dynamics**[1] that extend beyond any individual system and operate inside organizational and regulatory contexts. Privacy mechanisms, robustness techniques, and sustainability metrics provide the technical foundation; responsible AI adds bias detection, organizational governance structures, and stakeholder engagement processes on top of them. That integration is a significant infrastructure investment, and the training, inference, and memory overheads quantified in section 16.4.2 must be provisioned as essential capacity, analogous to how security and redundancy are budgeted in distributed systems.

Treating fairness, transparency, accountability, and privacy as rigorous engineering specifications rather than abstract ideals transforms responsible AI from aspiration into practice. Each principle becomes a design constraint with measurable criteria, attached to a specific stage of the machine learning lifecycle. The assurance lifecycle first establishes who may set and approve those constraints, and the deployment tier determines where each control can run and what it costs. Fairness, privacy, explainability, and safety are then developed one at a time, each from definition through measurement to the controls that enforce it. Accountability and sociotechnical analysis close the argument by testing whether those controls remain effective inside human workflows and institutions.

[1] **Sociotechnical system**: Coined by the Tavistock Institute in the 1950s to describe the interdependent relationship between humans and technology in the workplace. ML fleets are sociotechnical systems in exactly this sense, because their "performance" is not merely a benchmark score but an emergent property of how model outputs interact with user behavior, legal frameworks, and physical resource constraints.

16.2 The Governance and Assurance Lifecycle

A model registry can prove which weights reached production, but it cannot prove that the deployment was justified. That judgment requires a chain of evidence connecting the deployment context to affected people, foreseeable harms, measurable requirements, implemented controls, accountable owners, and the authority to launch or stop the system. Responsible AI therefore needs an assurance lifecycle: a repeatable operating path that converts normative commitments into reviewable claims and keeps those claims valid as the system changes.

The lifecycle begins before model selection. Teams first define the intended use, prohibited uses, affected parties, operating environment, and decision consequences. They then identify harms and requirements, including subgroup performance, privacy, security, robustness, sustainability, human-

↩ The model registry, which records each deployed version and its upstream and downstream dependencies, is established in section 12.3.1.

oversight, and contestability obligations. Some controls are implemented locally; others are inherited from data providers, foundation-model vendors, cloud platforms, or downstream operators. Each inherited control needs an interface contract stating what the supplier guarantees, what evidence supports that guarantee, and which residual risks remain with the deploying organization.

Constraint selection needs an explicit division of authority. Domain experts and affected parties identify consequential harms and acceptable uses; policy and legal specialists identify applicable obligations; engineers determine what can be measured and enforced. The accountable launch authority approves the resulting criteria and records unresolved disagreements. A convenient metric or an available implementation does not authorize the engineering team to choose whose errors matter most.

The resulting assurance case is not a single compliance document, but a versioned set of linked artifacts that structure accountability across the system lifecycle. At its foundation sits an explicit *claim* defining the property required for a specified use and operating envelope. An *argument* articulates how operational controls and empirical observations justify that claim, drawing on *evidence* from reproducible evaluations, audits, access logs, documentation, and user research. These claims and arguments hold only under stated *assumptions* regarding environmental, population, supplier, and workflow boundaries. To ensure operational authority, every assurance case specifies an accountable *owner* empowered to block launch, restrict use, or require remediation. An explicit *invalidation rule* governs that authority. It defines the measurable performance change, distribution shift, or operational incident that expires prior approval and triggers reassessment. Evidence reaches the claim only through the argument and only inside the stated assumptions, and the owner's launch decision lasts only until the invalidation rule fires (figure 16.1).

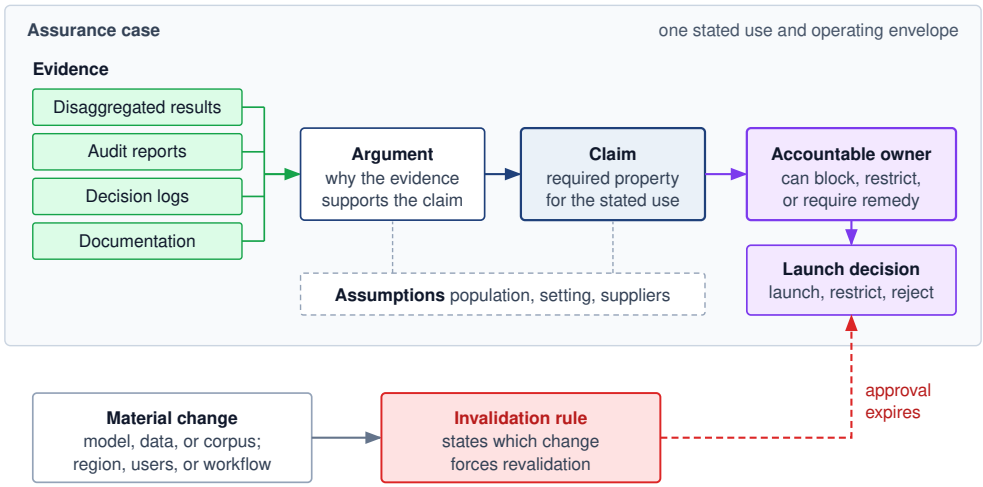

Figure 16.1: Assurance Case Structure: Evidence supports a claim only through an explicit argument and only within stated assumptions. An accountable owner turns the claim into a launch decision, and a material change named by the invalidation rule expires that approval, so the case is a versioned artifact rather than a one-time sign-off.

This structure prevents a common category error. A fairness dashboard, privacy accountant, red-team report, or carbon inventory is evidence for a bounded claim; none is itself proof that the full system is responsible. That evidence becomes operational only when an owner can act on it and affected people have a path to contest a decision and obtain remedy (table 16.1).

Table 16.1: Governance and Assurance Lifecycle: Responsible operation connects context and affected parties to requirements, inherited controls, evidence, launch authority, monitoring, incidents, remedy, material-change review, and retirement. Each stage produces an auditable artifact and a decision rather than a policy statement.

Lifecycle stage	Required artifact	Operational decision
Context and affected parties	Intended-use statement, prohibited uses, stakeholder map, decision-impact analysis	Is this use legitimate and within organizational authority?
Risk and requirements	Harm model, legal and policy obligations, measurable acceptance criteria	Which risks must be prevented, bounded, transferred, or accepted?
Controls and dependencies	Control inventory, supplier evidence, data and model lineage, output filters	Are local and inherited controls sufficient for the operating envelope?
Evaluation and assurance	Assurance case, disaggregated results with uncertainty, robustness and security evidence, sustainability accounting	Does the evidence support launch, restricted launch, or rejection?
Launch authority	Signed decision record, residual-risk owner, rollback and appeal paths	Who may authorize deployment, and who may stop it?
Operation and incident response	Monitoring thresholds, escalation owners, incident log, remedy workflow	When must the system degrade, roll back, notify, compensate, or suspend?
Material change	Change classification, impact assessment, revalidation scope	Does a new model, dataset, tool, region, or use invalidate prior approval?
Retirement	Decommission plan, retention and deletion record, user and dependent-system notice	How are access, data, models, interfaces, and unresolved harms closed out?

At release time, the evaluation contract defined in section 12.5.8 binds those assurance claims to the exact model, workload, acceptance authority, rollback target, and evidence-retention period. The assurance case explains why deployment is justified; the evaluation contract ensures that promotion and invalidation execute that judgment against one reproducible artifact.

Named frameworks organize this work at different levels. The National Institute of Standards and Technology (NIST) AI Risk Management Framework structures activity around Govern, Map, Measure, and Manage. ISO/IEC 42001 specifies an organizational AI management system, and ISO/IEC 23894 supplies AI risk-management guidance. The recommendations of the Organisation for Economic Co-operation and Development (OECD) and the United Nations Educational, Scientific and Cultural Organization (UNESCO) state broader principles and public responsibilities (Tabassi 2023; International Organization for Standardization and International Electrotechnical Commission 2023b, 2023a; Organisation for Economic Co-operation and Development 2024; UNESCO 2022). Mapped onto one operating lifecycle, those perspectives complement one another rather than serving as interchangeable checklists (table 16.2).

Table 16.2: Framework Crosswalk: Public frameworks differ in scope and legal status, but their durable systems contribution is complementary. They help organizations establish authority, identify and measure risk, operate controls, and reassess evidence through change and retirement.

Lifecycle responsibility	NIST AI RMF emphasis	Management-system and public-policy emphasis
Establish context, affected parties, and authority	Govern and Map	Organizational context, leadership, stakeholder participation, and human oversight
Measure risk and assemble evidence	Measure	Risk assessment, documented objectives, evaluation, transparency, and traceability
Decide, operate, and respond	Manage	Operational controls, accountability, incident response, remedy, and continual improvement
Reassess change and retirement	Govern and Manage	Change management, corrective action, lifecycle review, and decommissioning obligations

Regulation turns parts of this lifecycle into legal obligations, and it does so by classifying uses rather than by reducing AI risk to one verifiable metric. The EU AI Act (2024) uses a risk-based structure that includes prohibited practices, high-risk systems, transparency duties for some systems, and minimal or no-risk uses. High-risk deployments face conformity, documentation, and logging

obligations (European Parliament and Council of the European Union 2024). Narrower statutes govern particular uses, as the Illinois Artificial Intelligence Video Interview Act does for automated analysis of recorded job interviews (Illinois General Assembly 2020). United States federal policy has changed repeatedly since 2023. The safety reporting thresholds of Executive Order 14110 were revoked in January 2025, and later orders redirected agency emphasis toward AI innovation, cybersecurity, and critical-infrastructure defense (Executive Office of the President 2023, 2025a, 2025b, 2026). China's Interim Measures for Generative AI (2023) subject some public-facing services to security assessment and algorithm-filing requirements rather than a universal pre-release assessment (Cyberspace Administration of China et al. 2023).

These requirements cannot be collapsed into a single checklist. For a global ML fleet, applicable obligations depend on the service, affected population, organizational role, and jurisdiction, not server location alone. The deploying organization must identify the obligations that apply, assign owners to the procedures and evidence they require, and bind the approved policies to those deployment conditions in the control plane, the fleet layer that decides what may be deployed and where. The control plane preserves their versions and propagates changes to the components that enforce them.

↩ Governance as the control plane of the machine learning fleet is established in section 1.4.3.

Three details keep this lifecycle from becoming ceremonial. First, fairness estimates must carry denominators, confidence intervals, and data-coverage limits; a noisy subgroup gap should trigger investigation, not automatic claims of either discrimination or safety. Second, sustainability evidence belongs in the same assurance record, because the functional unit, system boundary, carbon signal, water accounting, embodied-impact allocation, uncertainty, and quality constraint determine whether an environmental claim is comparable or actionable. Third, a generative service that retrieves documents into a model's context needs an authorization boundary independent of model alignment. Least-privilege scopes, output validation, provenance, and revocation must be versioned and monitored because retrieved content remains an untrusted input.

↩ The sustainability measurement contract, which fixes the functional unit, system boundary, carbon signal, and attribution rule, is established in section 15.4.

Approval expires when its supporting assumptions no longer hold. A change of foundation model, training distribution, retrieval corpus, deployment region, user population, decision purpose, or human workflow can be *material* even when the application code is unchanged. The change record should identify which claims are invalidated and select the smallest defensible revalidation scope. At the end of service, retirement must revoke credentials, stop dependent interfaces safely, apply retention and deletion rules, preserve required audit evidence, notify affected users and downstream operators, and leave an owner for unresolved appeals or harms. Governance covers the system after launch and after shutdown, rather than solely at the moment a review board signs a release, and each stage of that assurance lifecycle depends on translating high-level principles into concrete checks across the ML pipeline.

16.3 Core Principles and the ML Lifecycle

An assurance case becomes enforceable only when its claims attach to specific pipeline gates: a continuous integration pipeline, for instance, can block a release when an evaluation detects a 15 percent drop in accuracy specifically for elderly users. Unlike a failing unit test or a latency regression, which the pipeline can check deterministically, a subgroup result requires a specified population, sufficient evaluation data, and an agreed acceptance criterion. Responsible AI translates those commitments into deployment gates. The continuous integration/continuous deployment (CI/CD) pipeline enforces the approved rule, while the accountable reviewer interprets uncertainty and decides whether the evidence supports release, restricted use, or remediation.

That release gate is one instance of a general translation from principle to enforceable requirement. Transparency discloses how the system as a whole was built and governed: its data sources, training procedure, intended uses, and known limitations. Explainability is narrower, producing evidence about the model's behavior on particular decisions for a specified audience. Together they make the system inspectable by the developers, auditors, and affected people who must judge it. Fairness constrains how selection rates, error rates, or the meaning of a score may differ across population subgroups; its violations can compound through feedback loops, as in predictive policing, while the service remains technically available. Privacy bounds what the system reveals about the people whose data it uses, safety and robustness bound how it fails, and accountability assigns owners who can trace, remedy, and answer for its outcomes.

Two further principles govern the objective and its supervision. Value alignment[2] asks whether the objective itself reflects legitimate human intent, and human oversight[3] preserves the ability of people to supervise, correct, or halt automated decisions. The sections that follow develop the first group in turn; alignment and oversight return where autonomy and human workflows make them binding, in section 16.8.3 and section 16.9.4.2.

No principle ensures a responsible system on its own. The principles overlap, compete for priority when budgets are finite, and act in different phases of the responsible AI lifecycle, from data collection through model training, evaluation, and deployment to monitoring (table 16.3).

Table 16.3: Responsible AI Lifecycle: Embedding fairness, explainability, privacy, accountability, and robustness throughout the ML system lifecycle, from data collection to monitoring, ensures these principles become architectural commitments rather than post-hoc considerations. The table maps these principles to specific development phases, revealing how proactive integration addresses potential risks and promotes trustworthy AI systems.

Principle	Data collection	Model training	Evaluation	Deployment	Monitoring
Fairness	Representative sampling	Bias-aware algorithms	Group-level metrics	Threshold adjustment	Subgroup performance
Explainability	Documentation standards	Interpretable architecture	Model behavior analysis	User-facing explanations	Explanation quality logs
Transparency	Data source tracking	Training documentation	Performance reporting	Model cards	Change tracking
Privacy	Consent mechanisms	Privacy-preserving methods	Privacy impact assessment	Secure deployment	Access audit logs
Accountability	Governance frameworks	Decision logging	Audit trail creation	Override mechanisms	Incident tracking
Robustness	Quality assurance	Robust training methods	Stress testing	Failure handling	Performance monitoring

Fairness and privacy constraints begin at data collection; robustness likewise depends on data quality and training choices, while accountability requires named owners before deployment. Their operational consequences become visible when the system serves users and teams must respond to failures. Explainability spans the full lifecycle, supporting model debugging at design time and evidence about decisions during use. This mapping reinforces that responsible AI is a multiphase architectural commitment, not a post-hoc compliance step, and how each phase is implemented depends on the deployment tier where the model runs.

16.4 Architectural Trade-Offs

A centralized audit can join predictions, demographic attributes, and delayed outcomes when access to those records is authorized. A federated model running on a million smartphones may provide no such joined dataset. The same fairness criterion then requires a different evidence path, such as restricted aggregate statistics or a separately governed audit sample. Deployment architecture determines where the evidence can be assembled and who can act on it.

These architectural differences redistribute both feasible safeguards and accountable parties. The comparison that follows uses four deployment tiers that differ in compute, connectivity, and update authority: *Cloud ML* runs on serving replicas in data-center fleets, whereas *Edge ML* moves execution to gateways and servers near the data source. On the client side, *Mobile ML* runs on phone-class devices with sustained power budgets of a few watts, the regime of Archetype C (Federated MobileNet). *TinyML* operates on microcontrollers with hundreds of kilobytes of memory and milliwatt power budgets, an extreme below Archetype C that often lacks any update channel. This book's other lighthouses sit in the cloud tier: Archetype A (GPT-4/Llama-3) and Archetype B (DLRM at Scale) serve from data-center fleets, and even Archetype C relies on a cloud aggregator for its fleet-level evidence. TinyML has no lighthouse of its own; it appears here because it marks the limit case, a deployment whose safeguards must be fixed before it ships.

Resource availability, latency constraints, user interface design, and the presence or absence of connectivity determine whether responsible AI principles can be enforced consistently across deployment contexts. The same model may support rich explanation, continuous monitoring, and

[2] **Value alignment**: The problem of ensuring AI systems optimize for human values rather than proxy objectives. Russell (2021) argues that specifying the objective is harder than optimizing it. The engineering consequence is that the training objective becomes a governed artifact with an owner and a review path, not a fixed input chosen once by the modeling team.

[3] **HITL (human-in-the-loop)**: Human participation in model decisions adds response delay and improves safety only when the operator can identify and correct the relevant error. Staffing demand depends on the decision arrival rate, the fraction escalated, and the time required per review. In ML serving architectures, HITL requires routing logic, escalation criteria, and queue management that reshape the inference pipeline.

↩ Archetype C (Federated MobileNet) and its few-watt per-device envelope are established in section 1.6.1.3.

centralized auditability in the cloud, yet require static validation and simplified safeguards when embedded in a device with intermittent connectivity. Tier choice is also an equity decision, because a safeguard that needs a data-center round trip reaches only users whose connectivity supports it (section 16.4.3).

16.4.1 Design trade-offs

Moving across those four tiers forces concrete trade-offs whenever two responsible AI goals compete for the same finite budget. An autonomous braking controller or a wearable gesture recognizer on the edge cannot spend hundreds of milliseconds computing feature attributions on every sensor frame, and a personalized model adapting on one phone cannot see whether its local updates widen error gaps across the rest of the population. Because each tier fixes how much compute, latency, connectivity, and data visibility a safeguard can draw, optimizing for one principle routinely constrains another.

Cloud-based systems can provision extensive monitoring, fairness audits, interpretability services, and privacy-preserving tools using shared compute and storage. Centralized data handling also concentrates exposure to surveillance and breaches. On-device execution can keep raw data local, but user control still depends on interfaces, permissions, retention rules, and update authority. Postdeployment visibility, fairness instrumentation, and adaptation depend on which evidence and maintenance paths the architecture provides.

Tensions between goals often become apparent at the architectural level. When an endpoint has too little latency budget for a costly explanation on every inference, designers can evaluate lightweight methods within that budget, precompute suitable summaries, or retain decision records for asynchronous analysis. The explanation's audience and required delivery time determine which of these paths can satisfy the obligation.

Conflicts also emerge between personalization and fairness. Systems that adapt to individuals based on local usage data often lack the global context necessary to assess disparities across population subgroups. To prevent personalized predictions from creating systematic exclusion, a system must balance user-level adaptation with mechanisms for group-level equity and auditability.

Privacy and robustness objectives can also conflict. Logging rare events or user outliers can help diagnose reliability failures, but the records may expose sensitive behavior. If raw records must remain local or encrypted, the improvement process needs an authorized evidence path, such as restricted review or aggregate diagnostics. Robustness requirements belong in the initial design; subsequent refinement depends on whether the retained evidence and update mechanism support it.

Energy budgets can constrain the frequency and placement of safeguards. Reducing model cost does not itself require weaker fairness monitoring; teams can compare local checks, sampled audits, and off-device analysis against the evidence and response deadlines. Privacy safeguards face the same constraint. A differentially private (DP) training mechanism clips each person's contribution and adds calibrated noise, so that adding or removing any one person's record changes the probability of any released result only within a stated bound. In federated learning, clipping, noise generation, and communication can consume additional energy, but their placement and cost depend on the mechanism. A central-DP design may add noise at the curator, a trusted server that combines the clipped updates, whereas a local-DP design places randomization on the client. Battery feasibility therefore requires measuring the chosen implementation.[4]

Together, these deployment examples reveal a broader systems-level challenge. Responsible AI principles cannot be considered in isolation. They interact, and optimizing for one may constrain another. The appropriate balance depends on deployment architecture, stakeholder priorities, domain-specific risks, the consequences of error, and the environmental and distributional impacts of computational resource requirements. Responsible machine learning design therefore depends on making those constraints explicit before choosing methods, so that the resulting system can be evaluated against the deployment context it must actually inhabit.

Comparing these settings (table 16.4) demonstrates that no deployment context dominates across all principles; each makes distinct architectural compromises. Cloud systems support computationally expensive explanation methods and centralized fairness monitoring. Shapley-value attribution (SHAP) credits each input feature with its average marginal contribution to a prediction across

[4] **Energy-privacy trade-off**: Privacy-preserving techniques like differential privacy and secure multi-party computation can increase computation, communication, and battery use. In federated learning on mobile devices, the equity implication is direct. Users with older devices, limited data plans, or limited battery life may be excluded from privacy-protected AI services, creating a system where privacy protection becomes contingent on hardware resources.

↩ Differential privacy, its privacy budget, and the accountant that tracks it are established in section 13.9.

subsets of the other features. Local interpretable model-agnostic explanations (LIME) fit a simple surrogate model to many perturbed copies of one input. Both require many extra model evaluations per explained decision, and section 16.7 develops them. Cloud deployments can afford that work, but introduce privacy risks through data aggregation. Edge and mobile deployments offer stronger data locality but limit postdeployment observability and global fairness assessment. TinyML systems face particularly tight budgets; their validation, monitoring, and update plans must reflect the connectivity and storage actually provided. These constraints also decide which users receive which safeguards, the equity question raised in section 16.4.3. Understanding these deployment constraints provides necessary context for the technical methods that translate responsible AI principles into production controls.

Table 16.4: Deployment Trade-Offs: Responsible AI principles manifest differently across deployment contexts due to varying constraints on compute, connectivity, and governance; cloud deployments support complex explainability methods, while TinyML severely limits them. Prioritizing certain principles like explainability, fairness, privacy, safety, and accountability requires careful consideration of these constraints when designing machine learning systems for cloud, edge, mobile, and TinyML environments.

Principle	Cloud ML	Edge ML	Mobile ML	TinyML
Explainability	Supports complex models and methods like SHAP and sampling approaches	Needs lightweight, low-latency methods like saliency maps	Requires interpretable outputs for users, often defers deeper analysis to the cloud	Severely limited due to constrained hardware; mostly static or compile-time only
Fairness	Large datasets allow bias detection and mitigation	Localized biases harder to detect but allows on-device adjustments	High personalization complicates group-level fairness tracking	Minimal data limits bias analysis and mitigation
Privacy	Centralized data at risk of breaches but can use strong encryption and differential privacy methods	Sensitive personal data on-device requires on-device protections	Tight coupling to user identity requires consent-aware design and local processing	Distributed data reduces centralized risks but poses challenges for anonymization
Safety	Vulnerable to hacking and large-scale attacks	Real-world interactions make reliability important	Operates under user supervision, but still requires graceful failure	Needs static safety envelopes, watchdogs, and fail-safe behavior
Accountability	Corporate policies and audits allow traceability and oversight	Fragmented supply chains complicate accountability	Requires clear user-facing disclosures and feedback paths	Traceability required across long, complex hardware chains
Governance	Applicable law and oversight; centralized evidence and named owners	Applicable law and oversight; allocate duties across integrators	Applicable law and oversight; coordinate platform and app duties	Applicable law and oversight; preserve evidence through servicing and retirement

The deployment analysis identifies where each obligation must be implemented. A constrained endpoint may log bounded decision metadata while a backend aggregates subgroup metrics, stores audit evidence, and serves explanations. An offline endpoint instead needs a defined evidence-retrieval and update path, or an operating envelope that does not depend on unavailable controls. Device capacity limits placement and frequency; it does not by itself remove the fleet operator's responsibility to provide the control.

Tier placement decides where a control runs; it does not decide which control a defect needs. As the hiring case at the start of the chapter showed, predictive accuracy on historical labels does not establish that the target is legitimate. Threshold postprocessing can change a decision-rate disparity, while an unsuitable target may require changes to labels and training, and both need validation and operational ownership.

Across the principle sections beginning in section 16.5, every technical control performs one of three jobs. Detection exposes a problem early enough to intervene, mitigation changes data, training, or decision policy to reduce the harm, and validation assembles the evidence a launch authority needs to accept the remaining risk. Deciding where in the four tiers each job can run requires first pricing what those mechanisms cost in accuracy, computation, and memory.

16.4.2 Computational overhead of responsible AI techniques

Adding adversarial training, which computes a seven-step attack on each batch before the training step, multiplies batch compute roughly eightfold, and running model-agnostic SHAP attributions on every request can multiply inference work by hundreds or thousands of evaluations. Detection, mitigation, validation, privacy, and explanation techniques each impose a distinct cost on training, inference, memory, and accuracy (table 16.5). Among SHAP algorithms, exhaustive enumeration evaluates the model on every subset of features, sampled KernelSHAP estimates the attributions from a random sample of subsets, and TreeSHAP computes them exactly from a tree model's structure (section 16.7.2). Pricing those overheads up front turns responsible AI methods from costless ethical labels into concrete capacity choices.

↩ Adversarial training is established in section 14.6.1.

Table 16.5: Performance Impact of Responsible AI Techniques: The table gives planning ranges for capacity exercises rather than universal benchmark results. Differential privacy and fairness constraints mainly affect training and validation budgets; explainability cost depends on the value function and algorithm: model-agnostic subset enumeration grows combinatorially, sampled kernel methods require repeated evaluations, and model-specific TreeSHAP can compute exact attributions in polynomial time under its tree and value-function assumptions. The differential-privacy accuracy range runs from fine-tuning a publicly pretrained model (low end) to training from scratch on a small image dataset (high end), the two regimes whose published private accuracies are compared in section 13.9.3. Adversarial training is shown as a multiple of standard training cost because each step first generates an attack; a seven-step attack costs about eight passes per batch. These sizing ranges help engineers reason about production constraints before choosing a responsible-AI implementation.

Technique	Accuracy loss	Training overhead	Inference cost	Memory overhead
Differential Privacy (DP-SGD)	1–20%	30–300%	Minimal	10–20%
Fairness-Aware Training	1–3%	5–15%	Minimal	5–10%
Sampled KernelSHAP	N/A	N/A	50–200%	20–100%
Exhaustive SHAP Enumeration	N/A	N/A	50–1,000×	50–200%
Adversarial Training	3–15%	2–11×	Minimal	50–100%
Federated Learning	5–15%	200–500%	Minimal	100–300%

These overhead ranges should be read as scenario assumptions, not as claims that any cited paper establishes a single universal overhead.[5] Actual overhead varies significantly based on model architecture, dataset size, attack strength, explanation method, hardware, communication pattern, and implementation quality. Responsible-AI controls consume real training, serving, memory, or review capacity, so they must be budgeted before launch. An organization unable to afford one configuration must assess alternatives such as a smaller model, sampled audits, or asynchronous explanations against the same required evidence. Resource limits can restrict the feasible deployment scope; they do not remove the obligation. Who can afford which configuration is therefore an equity question as well as a capacity question.

[5] **Measurement context**: The sizing exercise assumes common production reference points such as GPU training, GPU or CPU inference, image/text/tabular models, and benchmark datasets. These figures are not canonical constants; they are order-of-magnitude planning anchors for deciding whether a safeguard belongs inline, asynchronous, offline, or in the training pipeline.

16.4.3 Resource requirements and equity implications

A municipal clinic or school district operating on a tight computing budget cannot afford to double its training cluster for differential privacy or provision extra GPU replicas for real-time Shapley explanations. Each cell of table 16.3 is work that someone must pay for: representative sampling needs data acquisition, subgroup metrics need labeled evaluation data and analysis capacity, and user-facing explanations need serving capacity. When smaller organizations cannot fund every safeguard in the table, the people they serve receive fewer protections than the users of a better-resourced deployment. Large-scale deployments that depend on accelerator clusters and low-latency connectivity can likewise exclude users whose networks or devices cannot meet those assumptions.

The costs in table 16.5 set how high these barriers are, and they are paid in energy as well as capacity. Model-specific algorithms, approximation, caching, and selective generation change that

cost; asynchronous execution moves work off the request path without eliminating it. A high-traffic system serving these features to 10 million users must budget the work on the paths where it runs, then determine whether existing capacity is sufficient or additional resources are needed.

Infrastructure placement adds two further distributional effects. Data-center siting decisions that optimize power cost and latency can concentrate local environmental costs on nearby communities that had little say in the decision.[6] Connectivity also affects access to safeguards. A user-facing explanation that requires a low-latency round trip to a data center may miss its delivery deadline on a slow network, whereas server-side fairness monitoring can analyze recorded decisions without another user round trip. The architecture must distinguish immediate user needs from audit work that can run asynchronously. Computational equity therefore depends on placement decisions as well as algorithmic design, and the assurance case should record these distributional effects alongside the model's group metrics.

Computational costs and architectural constraints determine which controls a deployment can sustain. With those tier and capacity bounds in place, the next four sections examine each principle in turn, beginning with how a system defines, measures, and mitigates unfair disparities.

6 | **Data center environmental justice**: Data center siting can place backup diesel generators, cooling heat, and grid strain near communities that had little say in the siting decision; the distributional effect must be measured site by site rather than assumed.

16.5 Fairness and Bias Mitigation

A loan model can approve 70 percent of one group's applicants and 40 percent of another's while reporting a single headline accuracy. Exposing that gap, deciding which gap the deployment must close, and keeping it visible after launch are separate engineering tasks. Formal criteria define the disparity, statistical tests decide whether it exceeds sampling noise, and monitoring infrastructure tracks it in production.

16.5.1 Fairness in machine learning

Deciding whether that 30-point approval gap violates fairness requires more than comparing two averages, because automated decisions in hiring, lending, and healthcare inherit decades of structural inequality encoded in historical data. The first engineering task is to formalize "equitable treatment" as a conditional probability that a pipeline can measure and enforce.

Definition 16.2: Algorithmic fairness

Algorithmic fairness is a family of criteria for evaluating equitable treatment and outcomes in automated decisions. The group-based criteria developed here compare selection rates, errors, or score meaning across protected demographic groups; individual and causal criteria ask different questions.

1. **Significance**: Formal criteria make competing fairness goals measurable. Selection-rate or error-rate constraints can change predictive accuracy, while group-specific calibration asks whether predicted probabilities match observed outcome frequencies within each protected group. The trade-offs depend on the criterion, data, and decision policy. Satisfying one criterion does not establish equitable treatment by itself.
2. **Distinction**: Unlike average accuracy, which can hide subgroup disparities, group fairness examines the joint distribution of predictions and outcomes within each group, $p(\hat{Y}, Y \mid G)$. The chosen criterion determines which conditional rates to compare; a group gap diagnoses that metric rather than fairness as a whole.
3. **Common pitfall**: A frequent misconception is that there is a single "fair" solution. In reality, fairness requirements can conflict under specified assumptions, chiefly unequal base rates (the fractions of actual positive cases within groups), so choosing among them is a policy decision. Section 16.5.1.5 shows when and why the criteria defined below conflict.

Even when an evaluation disaggregates predictions by group, models trained on historical data can still reproduce and amplify systemic bias when the outcome label itself is a flawed proxy. An algorithm[7] used to allocate care management resources in US hospitals systematically underestimated the health needs of Black patients (Obermeyer et al. 2019). The model used healthcare

7 | **Healthcare algorithm scale**: The study examined one widely used algorithm in a broader class of risk-prediction tools that industry estimates placed at approximately 200 million U.S. patients annually (Obermeyer et al. 2019). In the study's simulated correction at the automatic-enrollment threshold (the risk-score cutoff above which patients were automatically identified for the program), the fraction of Black patients among those identified for additional care rose from 17.7 percent to 46.5 percent. These are shares of the selected population, not the fraction of all Black patients identified or a measured nationwide enrollment reduction.

expenditures as a proxy for health status, but due to longstanding disparities in access and spending, Black patients were less likely to incur high costs. As a result, the model inferred that they were less sick, despite often having equal or greater medical need. Seemingly neutral design choices such as proxy variable selection can therefore yield discriminatory outcomes when historical inequities are not accounted for. Correcting an invalid proxy can improve alignment with the actual clinical goal. A separate engineering question is whether imposing a particular fairness constraint reduces accuracy on a fixed target. That scenario-specific reduction, the fairness cost, is priced in section 16.5.3 once the criteria and the decision thresholds that enforce them have been defined.

Evaluating these risks requires formal criteria, defined in full below; the release gates of section 12.5.2.3 enforce two of them, demographic parity and equalized odds, as thresholds. Each criterion turns a different meaning of fairness (equal selection, equal error rates, equal opportunity for the qualified, or scores that mean the same thing in every group) into a conditional probability that a system can measure and enforce. Choosing a criterion is therefore choosing which of those meanings the deployment must honor. Suppose a model $h(x)$ predicts a binary outcome, such as loan repayment, and let G represent a sensitive attribute with subgroups a and b. Three widely adopted group criteria compare different conditional rates of that prediction.

16.5.1.1 Demographic parity

Definition 16.3: Demographic parity

Demographic parity is the fairness constraint where a model's positive prediction rate is independent of group membership ($\Pr(h(x) = 1 \mid G = a) = \Pr(h(x) = 1 \mid G = b)$).

1. **Significance**: It requires equal positive prediction rates across groups, regardless of underlying base-rate differences. It does not require identical outcomes for every individual, and it is not universally more restrictive than error-rate or calibration criteria.
2. **Distinction**: Unlike equalized odds (which focuses on error rates such as false positives), demographic parity focuses only on the final prediction, ignoring the relationship between the prediction and the ground truth.
3. **Common pitfall**: A frequent misconception is that demographic parity ensures "fairness." Equal approval rates can coexist with unequal error rates. Changing decision thresholds can alter who receives approval without changing the underlying scores or their calibration; the resulting decisions and the meaning of the scores require separate evaluation.

Demographic parity asks whether favorable outcomes, such as loan approval or treatment referral, occur at equal rates across subgroups defined by a sensitive attribute G. Formally, the model satisfies demographic parity if:

$$\Pr\left(h(x) = 1 \mid G = a\right) = \Pr\left(h(x) = 1 \mid G = b\right)$$

In the healthcare example, this criterion would ask whether Black and white patients were referred for care at the same rate, regardless of their underlying health needs. That may sound like equal access, but it ignores real differences in medical status and risk, potentially overcorrecting when needs are not evenly distributed. The limitation of ignoring base-rate differences motivates error-rate-based fairness criteria.

16.5.1.2 Equalized odds

Equalized odds requires that the model's predictions are conditionally independent of group membership given the true label. Specifically, the true positive and false positive rates must be equal across groups:

$$\Pr\left(h(x) = 1 \mid G = a, Y = v\right) = \Pr\left(h(x) = 1 \mid G = b, Y = v\right), \quad \text{for } v \in \{0, 1\}.$$

That is, for each true outcome $Y = v$, the model should produce the same prediction distribution across groups $G = a$ and $G = b$. Equalized odds equalizes conditional error rates, not error counts;

groups with different numbers of positive or negative cases can still experience different numbers of errors. Applied to a binary medical referral decision, this criterion requires equal referral rates across racial groups separately among patients who need the intervention and among those who do not, according to the chosen label Y.

16.5.1.3 Equality of opportunity

Equality of opportunity is a less stringent relaxation of equalized odds that focuses only on the true positive rate (Hardt et al. 2016). It requires that, among individuals who should receive a positive outcome, the probability of receiving one is equal across groups:

$$\Pr\left(h(x) = 1 \mid G = a, Y = 1\right) = \Pr\left(h(x) = 1 \mid G = b, Y = 1\right).$$

Equality of opportunity equalizes the approval rate among qualified individuals, who have $Y = 1$, across groups. It does not require identical decisions for every pair of qualified individuals or constrain errors among unqualified applicants.

In medical referral, this measure would require equal identification rates across racial groups among patients labeled as needing the intervention. Both error-rate criteria inherit whatever the label Y measures, which is the lesson of the healthcare case. A model trained to predict expenditure could equalize its error rates against expenditure and still misallocate care, because expenditure is not need. The cited study measured illness burden at a given risk score rather than group error rates, so it cannot itself be scored against these criteria. The loan scenario that follows supplies the binary labels and decisions needed to calculate all three on one model, and it becomes the section's running example.

Napkin Math 16.1: Calculating fairness metrics

Problem: Calculate demographic parity across two demographic groups (100 applicants each) for a loan approval model.

Variables:

- Group A: 70 approved (40 repaid, 30 defaulted), 30 rejected.
- Group B: 40 approved (30 repaid, 10 defaulted), 60 rejected.

Math:

1. **Group A approval rate**: $\Pr(h(x) = 1 \mid G = a) = \frac{70}{100} = 0.70$.
2. **Group B approval rate**: $\Pr(h(x) = 1 \mid G = b) = \frac{40}{100} = 0.40$.
3. **Parity gap**: $0.70 - 0.40 = 0.30$ (30 percentage-point gap).

Systems insight: Evaluating demographic parity exposes group approval rate disparities regardless of ground-truth repayment capability, requiring explicit metric selection during model threshold calibration.

The next calculation assumes that repayment labels are also known for rejected applicants. Among the rejected applicants in Group A, five would repay and 25 would default; in Group B, 20 would repay and 40 would default. These are hypothetical evaluation labels. A production lender ordinarily observes repayment only for approved loans, so estimating the same error rates requires additional evidence about rejected applicants and an account of how that evidence was obtained.

Napkin Math 16.2: Calculating equality of opportunity

Calculating equality of opportunity (true positive rate [TPR]):

Among applicants who *would actually repay* ($Y = 1$):

$$\Pr(h(x) = 1 \mid G = a, Y = 1) = \frac{40}{40+5} = \frac{40}{45} \approx 0.89$$
$$\Pr(h(x) = 1 \mid G = b, Y = 1) = \frac{30}{30+20} = \frac{30}{50} = 0.60$$

Disparity: $0.89 - 0.60 = 0.29$ (29 percentage-point gap in TPR)

The model violates equality of opportunity. Among qualified applicants who would repay, Group A members are correctly approved 89 percent of the time while Group B members are only approved 60 percent of the time.

Calculating equalized odds (true positive rate [TPR] and false positive rate [FPR]):

The TPR values are calculated in the preceding step. For false positive rates among applicants who *would not repay* ($Y = 0$):

$$\Pr(h(x) = 1 \mid G = a, Y = 0) = \frac{30}{30+25} = \frac{30}{55} \approx 0.55$$
$$\Pr(h(x) = 1 \mid G = b, Y = 0) = \frac{10}{10+40} = \frac{10}{50} = 0.20$$

The model also has unequal false positive rates. It incorrectly approves 55 percent of Group A applicants who will default, but only 20 percent of Group B applicants who will default. The model is therefore more lenient with Group A applicants who will not repay.

Systems insight: This model violates all three fairness criteria. Addressing one criterion does not automatically satisfy others. Whether they can all be satisfied together depends on base rates and prediction quality, the question the impossibility results below address.

16.5.1.4 Calibration

A fourth criterion concerns the score rather than the decision. A model satisfies score calibration with respect to a sensitive attribute if, among individuals assigned score s by the model, the fraction with positive outcomes equals that score within each group (Kleinberg et al. 2017):

$$\Pr(Y = 1 \mid s(x) = s, G = a) = s = \Pr(Y = 1 \mid s(x) = s, G = b), \quad \forall s$$

The condition applies at score levels supported within each group. Equal outcome frequencies across the groups are not enough; in each group, the outcome frequency at a score must equal that score. For hard binary predictions, a distinct condition is equal positive predictive value (PPV) (precision) among predicted positives. Groupwise score calibration checks each score level s separately, whereas PPV averages the outcome rate across all applicants whose scores clear the decision cutoff. Groupwise score calibration therefore does not by itself guarantee equal PPV after thresholding. Under calibration, a group's PPV equals the average score of its approved applicants, so if one group's approved applicants cluster just above the cutoff while the other group's sit higher in the score tail, the two PPVs differ. With a cutoff of 0.5, for example, a calibrated group whose approved applicants all score 0.6 has a PPV of 0.6, while a calibrated group whose approved applicants all score 0.8 has a PPV of 0.8. Equal PPV across groups requires

$$\mathrm{PPV}(a) = \frac{\Pr(Y = 1, h(x) = 1 \mid G = a)}{\Pr(h(x) = 1 \mid G = a)} = \mathrm{PPV}(b)$$

Evaluating this condition on the loan-approval example yields the respective predictive values:

$$\mathrm{PPV}(a) = \frac{40}{70} = 0.571$$
$$\mathrm{PPV}(b) = \frac{30}{40} = 0.750$$

The positive predictive value gap is $0.750 - 0.571 = 0.179$. Group B's predicted positives are actually positive 75 percent of the time, while Group A's are only 57 percent accurate. This violates predictive parity for hard decisions and reveals that the model is less reliable when predicting approval for Group A. Demonstrating score calibration itself would require comparing observed outcome frequencies with predicted probabilities within each group, using score bins with sufficient support. That check is essential in high-stakes decisions where people rely on predicted probabilities, because a miscalibrated model systematically over- or underpredicts risk for specific groups, misallocating resources and eroding trust.

16.5.1.5 Why the fairness criteria conflict

The loan model's failure to satisfy every criterion raises a broader question about their compatibility. Its answer depends on base rates and prediction quality, and the definitions can conflict outright under deployment-specific conditions.[8] A university admissions example makes the simplest conflict, between demographic parity and equalized odds, concrete before the broader calibration result (Kleinberg et al. 2017; Chouldechova 2017).

Suppose a university admits only qualified applicants, and a larger fraction q_a of Group A's applicants are qualified than the fraction q_b of Group B's. Equal opportunity asks that the same fraction r of qualified applicants be admitted from each group, for example 80 percent of qualified Group A applicants and 80 percent of qualified Group B applicants. The overall admission rates are then rq_a and rq_b, which differ whenever $q_a \neq q_b$, so demographic parity fails; they match only at $r = 0$.

The no-unqualified-admissions rule fixes each group's false positive rate at zero, so equal opportunity here coincides with equalized odds, and parity holds only when each group's true positive rate also falls to zero. Relaxing the no-unqualified-admissions rule removes the conflict between parity and equal opportunity, because admitting some unqualified Group B applicants can raise Group B's overall rate to Group A's. Equalized odds then fails, since Group B's false positive rate rises above Group A's zero.

The admissions case is one instance of a general condition. With unequal base rates, demographic parity and equalized odds hold together only if each group's true positive rate equals its false positive rate, so a positive decision is equally likely among actual positives and actual negatives and carries no information about the outcome. A constant classifier satisfies both; any informative classifier must give one up. Which constraint to give up is a value judgment about the decision, not a property the model can settle.

The fairness impossibility theorem (principle 26) is the canonical result, and it concerns scores rather than decisions. It joins the groupwise calibration of section 16.5.1.4 to score balance, which requires equal average scores across groups among people with positive outcomes and, separately, among people with negative outcomes. With unequal base rates and imperfect prediction, all three requirements cannot hold together (Kleinberg et al. 2017) (figure 16.2). The admissions conflict concerns decisions and this result concerns scores, so they constrain different criteria. The score result also does not decide whether a calibrated score meets equalized odds at a particular threshold; that narrower question is answered by computing the decision-level error rates at that threshold, as the loan example did.

Satisfying one criterion may preclude another, so engineers must treat fairness metrics like latency budgets, explicit trade-offs chosen by stakeholders, enforced by the system, and monitored for violation. Determining which metric to prioritize requires careful consideration of the application context, potential harms, and stakeholder values as detailed in section 16.9.4.3.

The conflict persists as the classification threshold moves. In the synthetic two-group example of figure 16.3, where the groups have unequal base rates, every threshold from 0.15 to 0.85 leaves at least one of demographic parity, equalized odds, and equal opportunity substantially violated. The exact pattern depends on the chosen score distributions. A threshold that reduces one measured gap need not minimize the others, so the policy must specify which criterion governs the decision and evaluate the remaining gaps.

Group metrics also leave open how an observed disparity arose, because a gap does not identify the mechanism that produced it. Causal analysis addresses that question through explicit assumptions.

[8] **Fairness impossibility theorems**: Kleinberg et al. (2017) studies groupwise score calibration together with balance of mean scores in each outcome class. Chouldechova (2017) studies predictive parity and error-rate constraints for recidivism prediction. These are distinct conditions; their incompatibility results require assumptions on base rates and predictive performance. Demographic parity is a distinct constraint on selection rates that can also conflict with calibration or error-rate goals in practice. The systems consequence is that no amount of engineering can satisfy every normative fairness criterion in all settings, so fairness becomes a constrained multi-objective optimization requiring explicit policy choices about which criterion to prioritize for a given deployment context.

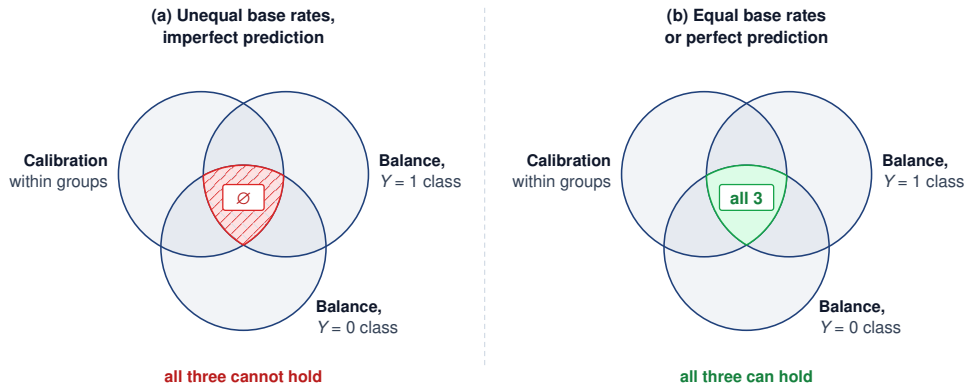

Figure 16.2: Conditional Fairness Trade-Offs: For score-based prediction, calibration within groups and equal mean scores within both outcome classes cannot all hold with unequal base rates and imperfect prediction. Equal base rates or perfect prediction are exceptional cases. This score-balance theorem is distinct from checking equalized odds at one decision threshold.

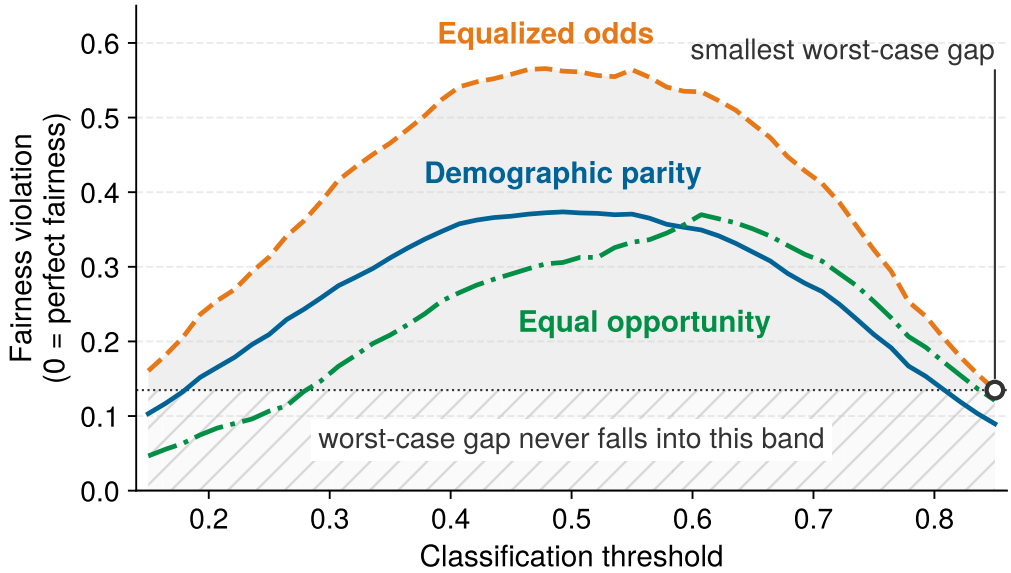

Figure 16.3: Fairness Metric Disagreement Across Thresholds: Three fairness metrics, demographic parity, equalized odds, and equal opportunity, computed on synthetic scores for two groups with differing base rates as the classification threshold varies over a nontrivial operating range. At every evaluated threshold, at least one metric is substantially violated, illustrating practical trade-offs among selection-rate, error-rate, and opportunity criteria. Trivial always-positive or always-negative classifiers fall outside the evaluated range. The shaded region marks the largest of the three violations at each threshold.

16.5.1.6 Observational and causal fairness

The preceding metrics describe associations in observed decisions, such as group differences in selection rates, error rates, or calibration. The same approval-rate gap could arise from direct use of a protected attribute, a correlated proxy, unequal access to an earlier process, measurement error, or an unobserved common cause. No comparison of observed rates alone distinguishes these mechanisms.

A causal claim adds a structural causal model whose arrows state how protected attribute G, mediators M (variables on a path from group membership to the decision), outcome Y, and decision $\hat{Y}$ are assumed to relate. **Counterfactual fairness** then asks whether the decision for the same individual, holding fixed the modeled background factors (the person's characteristics that group membership does not cause), would have changed in a counterfactual world where G had taken

another value (Kusner et al. 2017). This is a question about an assumed data-generating process, not a stronger version of demographic parity.

The healthcare proxy illustrates the distinction. A graph might include race $G \to$ access to care $M_1 \to$ expenditure $M_2 \to$ risk score $\hat{Y}$, alongside pathways from health need Y to expenditure. A total effect asks how changing G propagates through every modeled path. A path-specific effect isolates selected routes, such as the historically mediated access-to-expenditure path (Nabi and Shpitser 2018). Declaring one path permissible and another discriminatory is a domain and normative judgment; neither the graph nor the estimator makes that choice automatically.

Causal conclusions require assumptions that observational data alone cannot establish: which variables and arrows belong in the graph, whether relevant confounders (common causes of two modeled variables) are measured, and whether the counterfactual interventions are coherent. Different plausible structures can support different answers. Causal analysis therefore complements group metrics, domain knowledge, legal analysis, and stakeholder participation. The metrics continue to monitor realized disparities; the causal model makes the additional mechanism claim and its assumptions inspectable.

16.5.2 Measuring disparity

Even before a team models causal pathways, monitoring realized disparities in production faces a statistical challenge: a 30-point approval gap observed on 20 loan applications could be sampling noise, and a single-axis audit that checks race and gender separately can report high accuracy on both while masking severe errors where those attributes intersect. Turning a fairness criterion into evidence a launch authority can act on requires a point metric with a denominator and a confidence interval, coverage of the intersections where errors concentrate, and a governed choice of the criterion the deployment will be judged against.

16.5.2.1 Quantitative fairness measurement

Making that evidence operational starts with reducing a conditional-probability gap to a scalar summary that an audit or release gate can threshold. The mathematical framework quantifies the magnitude of the disparity and its uncertainty, while domain and policy owners must justify the acceptance criteria and authorize the response. A computed gap supplies evidence for that decision rather than deciding which harms are acceptable.

Four point metrics turn the fairness criteria introduced above into numbers an auditor can threshold. Each one equalizes a different quantity, so each exposes a different failure of the same classifier. On the running loan example (table 16.6), Group A is approved at 0.70 and Group B at 0.40. Qualified-applicant true-positive rates are 0.89 and 0.60, and unqualified-applicant false-positive rates are 0.55 and 0.20.

Table 16.6: Fairness Metrics on the Loan Example: Four group-fairness metrics, each equalizing a different quantity, evaluated on the loan approval model. The disparate impact ratio is a descriptive selection-rate screen whose legal weight depends on the application; the additive differences are easier to track over time; the average absolute odds difference captures both true-positive and false-positive error gaps.

Metric	Equalizes	Formula (a, b)	Loan value	Choose it when
Disparate impact (DI) ratio (four-fifths rule) (Feldman et al. 2015)	lower approval rate divided by higher approval rate	$\frac{\min_g \Pr(h=1 \mid g)}{\max_g \Pr(h=1 \mid g)}$	$\frac{0.40}{0.70} = 0.57$	a descriptive selection-rate screen is useful; legal interpretation depends on the application
Statistical parity difference (SPD) (Calders and Verwer 2010)	additive gap in approval rates	$\Pr(h=1 \mid a) - \Pr(h=1 \mid b)$	$0.70 - 0.40 = 0.30$	comparing across groups or tracking drift against a deployment-approved gap threshold

Metric	Equalizes	Formula (a, b)	Loan value	Choose it when
Equal opportunity difference (Hardt et al. 2016)	true-positive rates among qualified applicants	$\Pr(h=1 \mid a, Y{=}1) - \Pr(h=1 \mid b, Y{=}1)$	$0.89 - 0.60 = 0.29$	the harm is denying deserving individuals
Average absolute odds difference (Hardt et al. 2016; Bellamy et al. 2019; Bird et al. 2020)	both true-positive and false-positive rates	$\frac{1}{2}\left[\lvert\Delta\text{TPR}\rvert + \lvert\Delta\text{FPR}\rvert\right]$	$\frac{1}{2}[0.29 + 0.35] = 0.32$	both error types matter; the metric is 0 at perfect equalized odds

The four metrics measure different quantities by construction. The disparate impact ratio of 0.57 means the lower-approval group receives approvals at only that fraction of the higher-approval group's rate, below the illustrative 0.8 screening threshold. That threshold comes from the four-fifths rule, a U.S. employment-selection convention that flags any group whose selection rate falls below 80 percent of the highest group's rate. The additive 0.30 gap, the 0.29 opportunity gap, and the 0.32 average absolute odds gap each quantify a distinct dimension of the same disparity. No single number is sufficient, which is why the impossibility results of section 16.5.1.5 force an explicit choice about which metric a deployment will prioritize.

16.5.2.2 Statistical testing of fairness violations

Confidence intervals and hypothesis tests help assess whether sampling variation plausibly explains an observed disparity under a stated sampling model. A test starts from a null hypothesis, the assumption that the groups share one true rate so that any observed gap is sampling noise, and asks how unlikely the observed gap would be if that assumption held. For demographic parity, the null hypothesis $H_0 : \Pr(h(x) = 1 \mid G = a) = \Pr(h(x) = 1 \mid G = b)$ can be tested using a two-proportion z-test when observations are independent and expected approval and rejection counts in both groups support the normal approximation. Repeated decisions for the same person or correlated observations require an uncertainty estimate that accounts for that dependence. Under the independent-sample assumptions, the test statistic is:

$$z = \frac{\hat{p}_a - \hat{p}_b}{\sqrt{\hat{p}(1-\hat{p})\left(\frac{1}{n_a} + \frac{1}{n_b}\right)}}$$

where $\hat{p}_a$ and $\hat{p}_b$ are the sample approval rates, $\hat{p} = \frac{n_a\hat{p}_a + n_b\hat{p}_b}{n_a + n_b}$ is the pooled proportion, and n_a, n_b are sample sizes.

For the loan example with $n_a = n_b = 100$, $\hat{p}_a = 0.70$, $\hat{p}_b = 0.40$:

$$\hat{p} = \frac{100(0.70) + 100(0.40)}{200} = 0.55$$
$$z = \frac{0.70 - 0.40}{\sqrt{0.55(0.45)(0.02)}} = \frac{0.30}{0.0704} = 4.26$$

With $z = 4.26$ (far exceeding critical value $z_{\alpha/2} = 1.96$, the value the statistic exceeds in magnitude only 5 percent of the time when the null hypothesis holds), the test rejects H_0 and concludes the demographic parity violation is statistically significant at $p_{\text{val}} < 0.001$, where the p-value p_{val} is the probability of a gap at least this large under H_0. A confidence interval reports the same evidence as a range of plausible gaps, the observed gap plus or minus the critical value times the gap's standard error, and it narrows only with the square root of the group sample sizes. Similar tests can be constructed for equal opportunity and equalized odds by restricting to subpopulations where $Y = 1$ or $Y = 0$ respectively. Statistical significance does not establish the severity or acceptability of a disparity. Even a small rate difference can affect many people at fleet scale; the decision depends on the consequences and approved criteria as well as magnitude. Conversely, large disparities in small samples may not reach statistical significance but still warrant investigation or intervention. A significant gap also says nothing about its cause, the question that section 16.5.1.6 addresses.

16.5.2.3 Intersectional fairness

Even a statistically significant test on race or gender alone can miss where errors concentrate: in the Gender Shades evaluation, which tested commercial gender classifiers on face images grouped by skin type and gender, IBM's gender-classification system achieved 99.7 percent accuracy for lighter-skinned men but only 65.3 percent for darker-skinned women (Buolamwini and Gebru 2018). Standard fairness analysis that evaluates single axes of identity independently obscures such compounded disparities, a failure mode sometimes called fairness gerrymandering, in which a model looks fair on each attribute taken alone while failing subgroups defined by their combination. Detecting and mitigating these gaps requires evaluating model performance on intersectional subgroups (for example, Race×Gender) rather than marginal averages alone.

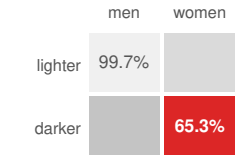

Error concentrates at the intersection: darker-skinned women.

Errors fell after the disparity was published, although the evidence shows an association rather than a controlled effect. In August 2018, about six months after the Gender Shades results were published, Raji and Buolamwini (2019) repeated the audit on the same benchmark. Each of the three companies that had been named had reduced its error on darker-skinned women: IBM from 34.7 to 17.0 percent, Face++ from 34.5 to 4.1, and Microsoft from 20.8 to 1.5 (figure 16.4). Two systems the original audit had not named, from Amazon and Kairos, measured for the first time in 2018, misclassified 31.4 and 22.5 percent of darker-skinned women, comparable to the named systems' 2017 levels. The result still suggests that a specific, published intersectional audit can act as an operational lever as well as a measurement.

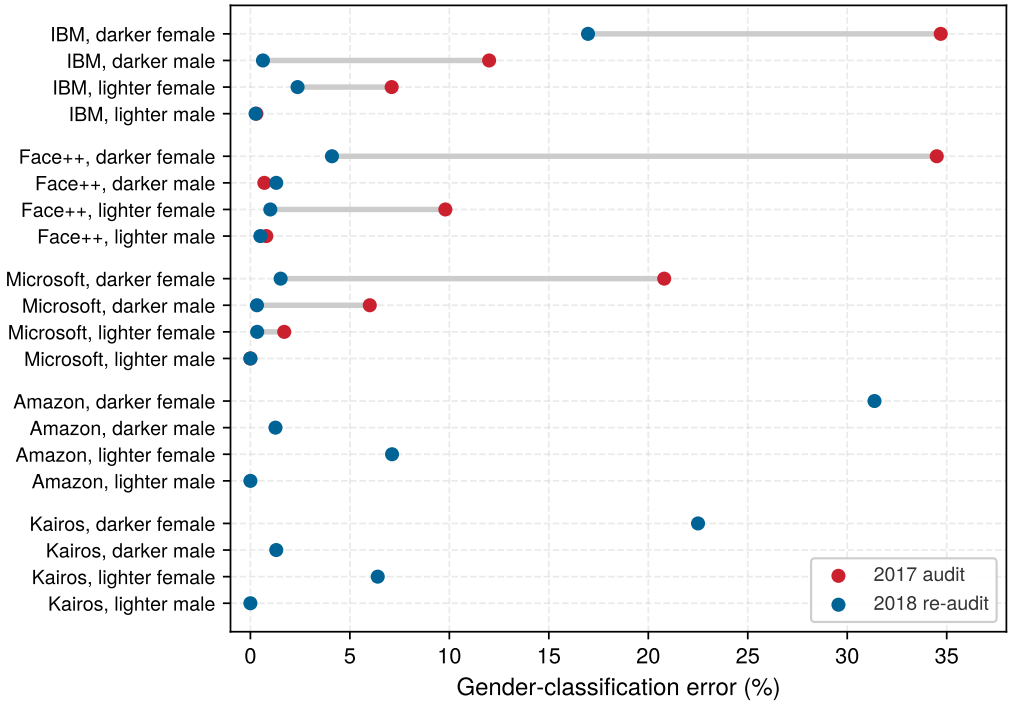

Figure 16.4: Audit and Re-audit: Gender-classification error by skin-type and gender subgroup on the Pilot Parliaments Benchmark, for the three systems named in the 2018 Gender Shades audit (measured in spring 2017, red) and in an August 2018 re-audit (blue), and for two systems first measured in the re-audit (blue only). Errors fell most for the darker-skinned women the audit had highlighted. Data from Buolamwini and Gebru (2018) (Table 4) and Raji and Buolamwini (2019) (Table 1).

An intersectional audit must specify the population slices it can identify and the evidence available within them. Work on multicalibration extends calibration checks to a specified, potentially rich family of overlapping subgroups (Hébert-Johnson et al. 2018). This does not certify every possible subgroup or every fairness criterion. It motivates evaluation and monitoring infrastructure that can examine relevant intersections while reporting where sample support or attribute availability limits the claim.

16.5.2.4 Choosing the governing criterion

The running loan example establishes how to measure a disparity and its sampling uncertainty. Applying that framework still requires identifying whether a positive decision grants a benefit or imposes a burden, checking the criteria a headline metric misses, and choosing the requirement that fits the deployment.

Napkin Math 16.3: Auditing a confusion matrix

Problem: A fraud detection model operates on two groups. What do demographic parity and equal opportunity reveal about the deployment?

Variables:

- **Group A (majority)**: TP $= 450$, FP $= 50$, FN $= 30$, TN $= 470$ ($n_{\text{records}} = 1000$).
- **Group B (minority)**: TP $= 180$, FP $= 70$, FN $= 120$, TN $= 630$ ($n_{\text{records}} = 1000$).

Math:

1. **Demographic parity (positive prediction rate)**: $\Pr(\hat{Y} = 1)$.
   - Group A: $(450 + 50)/1000 = 0.50$.
   - Group B: $(180 + 70)/1000 = 0.25$.
   - **Gap**: 0.25. The lower-to-higher flag-rate ratio is $0.25/0.50 = 0.5 < 0.8$. Fraud flags are adverse outcomes, so this descriptive ratio is not a four-fifths employment-selection test.
2. **Equal opportunity (TPR)**: $\text{TP}/(\text{TP} + \text{FN})$.
   - Group A: $450/(450 + 30) = 0.938$.
   - Group B: $180/(180 + 120) = 0.60$.
   - **Gap**: 0.338. (Severe violation).

Systems insight: Here a positive prediction flags suspected fraud, rather than granting a benefit as in the loan example. Equalizing detection rates therefore requires weighing missed fraud against false accusations. Lowering Group B's threshold may catch more fraud but can also increase false alarms; the score distributions are needed to quantify the effect on predictive parity.

The running loan model fails every criterion at once, which can suggest that a single headline metric would have caught it. A second hypothetical loan model, deliberately different from the running example, shows the opposite risk, in which one criterion is satisfied exactly while the others fail (table 16.7). Unlike the running model, whose false positive rates differ (0.55 against 0.20), this model approves would-be defaulters at the same 20 percent rate in both groups. Equalized odds would require both true positive rates and false positive rates to match, so the table shows the narrower case of false-positive-rate parity rather than full equalized odds. Because one fairness metric can appear satisfied while others expose substantial disparities, a deployment must decide in advance which criterion governs it.

Table 16.7: Disaggregated Fairness Metrics: A hypothetical loan approval system satisfies equalized false positive rates (0 percentage-point gap) but violates demographic parity (15 percentage-point approval gap) and equal opportunity (30 percentage-point TPR gap). Production systems must document the selected criteria and evaluate candidate policies against them.

Metric	Definition	A	B	Gap
Approval Rate	(TP + FP) divided by total	55%	40%	15 pp
True Positive Rate	TP divided by positives	90%	60%	30 pp
False Positive Rate	FP divided by negatives	20%	20%	0 pp

Metric	Definition	A	B	Gap
Positive Predictive Value	TP divided by predicted positives	82%	75%	7 pp

Example 16.1: Choosing a fairness metric

Scenario: A hiring recommendation engine must select its primary fairness constraint before production launch across heterogeneous applicant pools. The accountable reviewers prioritize equal access for qualified applicants and have independently validated qualification labels for both selected and rejected applicants.

Diagnosis: Demographic parity forces equal selection rates regardless of candidate qualification, while calibration requires predicted probabilities to match outcome frequencies within each group. Equal opportunity specifically mandates equal true positive rates among qualified applicants ($\text{TPR}_a = \text{TPR}_b$), directly matching the goal of equal opportunity for qualified talent.

Systems lesson: Metric selection is a governance decision that sets the operational objective of the serving pipeline. System designers must align the mathematical constraint with the domain's legal and ethical requirements before training or thresholding.

The chosen fairness definition determines which error trade-off the system is allowed to make, so the choice precedes optimization against that criterion. A model can satisfy a specified fairness criterion, but responsibility depends on the system that uses it. The system boundary extends past the model's weights to the pipeline and infrastructure that run them, because who can reach the safeguards and who bears their environmental costs are distributional effects of the same system (section 16.4.3). The fairness audit itself is part of that system, and the demographic attributes it collects can increase privacy exposure.

Governance must therefore justify both the fairness claim and the data practices used to support it across the full lifecycle (table 16.3). Data collection, labels, modeling choices, evaluation procedures, and decision thresholds all shape the resulting outcomes, and postprocessing can repair a selected rate disparity without resolving an invalid target or an unsupported deployment. When the target is valid and the model is already trained, adjusting the classification threshold is the most direct lever for trading off competing group rates and overall accuracy.

16.5.3 Threshold setting and fairness trade-offs

When two applicant groups have different score distributions, applying a single global approval cutoff admits different fractions of qualified applicants from each group, whereas lowering one group's cutoff to equalize opportunity gives two applicants with the same score different decisions. Postprocessing changes fairness metrics after training by adjusting classification thresholds per group rather than retraining the weights. In the equations below, a larger score favors a positive decision, as with predicted repayment probability. Given a scoring function $s_{\text{dec}}(x)$, define group-specific thresholds $\tau_{\text{thr},a}$ and $\tau_{\text{thr},b}$ such that $h_a(x) = \mathbb{1}[s_{\text{dec}}(x) \geq \tau_{\text{thr},a}]$ for group a and similarly for group b.

To achieve demographic parity, solve:

$$\Pr(s_{\text{dec}}(x) \geq \tau_{\text{thr},a} \mid G = a) = \Pr(s_{\text{dec}}(x) \geq \tau_{\text{thr},b} \mid G = b)$$

To achieve equal opportunity, solve:

$$\Pr(s_{\text{dec}}(x) \geq \tau_{\text{thr},a} \mid G = a, Y = 1) = \Pr(s_{\text{dec}}(x) \geq \tau_{\text{thr},b} \mid G = b, Y = 1)$$

For equalized odds, both true positive and false positive rate constraints must hold simultaneously. Varying a single deterministic threshold for a group moves its true positive rate and false positive rate together along one receiver operating characteristic (ROC) curve, so when two groups have different ROC curves, no single cutoff per group may land on the same (FPR, TPR) pair. Postprocessing can instead randomize between two threshold rules within each group, which blends their error rates

linearly and reaches interior (FPR, TPR) points where the two groups' achievable regions overlap (Hardt et al. 2016). The resulting criterion constrains group-level probabilities, so validation must evaluate the randomized policy rather than only its component thresholds.

Threshold adjustment still requires evaluating the other criteria. When base rates differ (that is, $\Pr(Y = 1 \mid G = a) \neq \Pr(Y = 1 \mid G = b)$), calibration and error-rate constraints can conflict under the conditions discussed earlier. The effect of a particular threshold policy must be measured on the joint score and label distribution; a base-rate difference alone does not determine its accuracy cost.

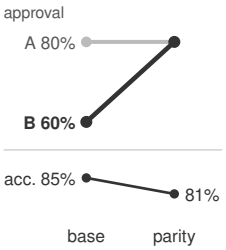

Lifting Group B's approval rate to parity costs 4 points of accuracy.

The accuracy cost of a threshold policy becomes concrete on a separate credit model, distinct from the running loan example. Its score is a predicted default probability, so a lower score favors approval and each cutoff admits applicants below it. The scenario enforces demographic parity, the first threshold equation above, and prices the accuracy that enforcement costs.

Napkin Math 16.4: The fairness cost

Problem: Consider a credit model with 85 percent accuracy. Group A (majority) has a 20 percent default rate. Group B (minority) has a 40 percent default rate due to systemic factors. If demographic parity (equal approval rates) is enforced, what happens to accuracy?

Math:

This scenario holds Group A's approval rate fixed and raises Group B's rate to match it; demographic parity alone does not require that particular choice. The accuracy values are assumed held-out validation results, where the additional Group B approvals carry higher default risk.

1. **Unconstrained**: Model approves everyone with predicted default prob < 30 percent.
   - Group A approval: 80 percent.
   - Group B approval: 60 percent.
   - Total Accuracy: 85 percent.
2. **Constrained (parity)**: Holding Group A's rate fixed, approve Group B at 80 percent rate.
   - New threshold for Group B: Approve default prob < 50 percent.
   - The revised cutoff admits Group B applicants with higher predicted default risk than the original policy allowed.
   - New Total Accuracy: 81 percent.

Systems insight: This parity policy cost 4 percentage points of accuracy (about a 4.7 percent relative accuracy drop in this credit-scoring scenario). The drop is the fairness cost of this scenario, not a universal property of fairness constraints. Accuracy still measures prediction against the specified default labels; whether those labels capture the legitimate lending objective is a separate question. Engineers must report the predictive trade-off while making the target and policy assumptions explicit.

The 4 percentage points cost belongs to one parity policy on one prediction target. A different criterion moves the thresholds for a different reason, and the next scenario, which returns to the repayment-score convention, targets equal opportunity by comparing the score distributions among qualified applicants. Differential thresholds require access to sensitive attributes at inference time and raise concerns about explicit group-based treatment, which may itself be considered unfair or illegal in certain jurisdictions. The required thresholds depend on those distributions and the chosen criterion; a lower threshold for the minority group is not a universal fairness rule.

Napkin Math 16.5: Threshold for equal opportunity

Problem: Calculate group-specific thresholds required to equalize true positive rates (equal opportunity) across Group A and Group B in a credit scoring model.

Variables: Let the score $s_{\text{dec}}(x)$ be an unbounded decision score, not a calibrated probability, and write Group A as $G = a$ and Group B as $G = b$. Assume $s_{\text{dec}}(x) \mid (Y = 1, G = a) \sim \mathcal{N}(0.72, 0.15^2)$ and $s_{\text{dec}}(x) \mid (Y = 1, G = b) \sim \mathcal{N}(0.65, 0.16^2)$. These are hypothetical positive-class distributions; no negative-class distribution or prevalence is specified.

Math:

1. **Single threshold**: At $\tau_{\text{thr},a} = \tau_{\text{thr},b} = 0.60$, $\text{TPR}_a = 1 - \Phi(-0.8) \approx 0.788$ and $\text{TPR}_b = 1 - \Phi(-0.3125) \approx 0.623$, where Φ is the standard normal cumulative distribution function.
2. **Equalizing thresholds**: Preserve Group A's threshold and match its standardized cutoff, $(\tau_{\text{thr},b} - 0.65)/0.16 = -0.8$. Thus $\tau_{\text{thr},b} = 0.522$ yields the same true positive rate.
3. **Unspecified costs**: Lowering Group B's threshold cannot decrease its false positive rate, but its magnitude requires $s_{\text{dec}}(x) \mid (Y = 0, G = b)$. Precision also requires Group B's prevalence. Neither quantity follows from the positive-class distributions alone.

Systems insight: Equal opportunity determines a threshold only relative to the stated score distributions and target rate. Before deployment, evaluate the resulting false positives, precision, and review workload on representative data. Threshold policies change decisions, not the underlying score calibration.

A separate illustrative loan scenario partitions applicants into Subgroup A (higher score distribution, blue band) and Subgroup B (lower score distribution, amber band), and the summary panels of figure 16.5 contrast two operating regimes. Applying a single global threshold ($\gamma_{\text{cls}} = 0.75$) approves qualified applicants in Subgroup B at only 25 percent TPR compared to 75 percent for Subgroup A, a threefold opportunity disparity. Conversely, shifting to group-adapted thresholds ($\tau_{\text{thr},a} = 0.75$ for Subgroup A and $\tau_{\text{thr},b} = 0.45$ for Subgroup B) equalizes true positive rates at 75 percent. That policy requires the authorized decision component to access the relevant group attribute and log the governing threshold. Neither the attribute lookup nor the policy decision must run on the client, but the system must preserve evidence explaining why identical scores receive different decisions.

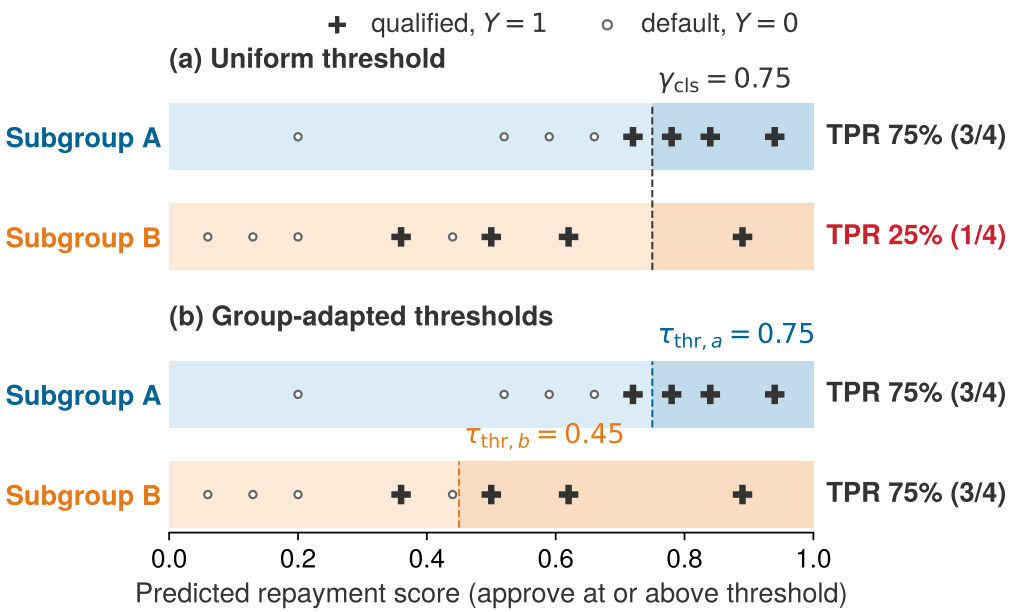

Figure 16.5: Threshold-Dependent Fairness: Varying classification thresholds across subgroups allows equal true positive rates but introduces complexity in model serving and requires access to sensitive attributes at inference time. Achieving fairness requires careful consideration of subgroup-specific performance, as a single threshold may disproportionately impact certain groups, highlighting the tension between accuracy and equitable outcomes in machine learning systems.

16.5.4 Fairness evidence across deployment tiers

Group-adapted thresholds and disaggregated audits assume that an authorized component can join predictions, demographic attributes, and outcomes, an assumption that breaks when execution leaves the data center. A federated keyboard model running across millions of phones, in the same federated regime as Archetype C (Federated MobileNet), never sends raw keystrokes or user demographic labels to the central server, so the fleet operator cannot run a central SQL query to check whether next-word accuracy drops on a protected subgroup. Fairness then becomes a deployment problem whose binding constraint is data visibility. A cloud platform with authorized access to demographic labels and outcomes can compute group-level metrics, run fairness-aware training, and audit post-hoc. A federated[9] or on-device deployment instead needs an explicit evidence path, such as aggregate client statistics or a separately governed audit sample. Without sufficient evidence, the corresponding metric is unmeasurable; local execution alone does not make the audit impossible. The evidence available for fairness analysis narrows from cloud to TinyML (fairness row of table 16.4), and that narrowing makes two recurring problems harder to check.

[9] **Federated learning and fairness**: Google's public Gboard trial illustrated training across devices without centralizing raw examples (Hard et al. 2018). Model updates alone do not necessarily reveal the demographic labels and outcomes needed for a fairness audit, so the communication cost, privacy protection, and statistical uncertainty of whatever evidence path replaces them must be evaluated for that design.

The first problem concerns decision policy. Equal aggregate accuracy across groups does not establish equal opportunity under a uniform threshold; in the illustration of figure 16.5, one shared threshold leaves the two groups' true positive rates far apart. A deployment must select and evaluate a policy against its approved criterion; group-specific thresholds are one possible intervention, not a requirement imposed by edge execution. The second problem sharpens the personalization conflict raised in section 16.4.1. Local adaptation aims to improve a user's utility, whereas individual fairness requires similar individuals to receive similar decisions under a task-relevant similarity metric (Dwork et al. 2012). Personalization does not establish that requirement, and adaptation on sparse or unrepresentative observations can reinforce existing gaps. When the relevant evidence stays local, fleet operators need an authorized way to assess the resulting disparities. Section 16.9.4.1 analyzes the feedback mechanism and its intervention points; the deployment consequence is that a fleet which cannot centralize outcomes must budget an explicit evidence path for detecting gaps that personalization reinforces.

That evidence path creates a further responsible AI tension. Keeping sensitive data local reduces centralized exposure, while fleet-level audits still need adequate group statistics. Privacy architecture (section 16.6.2) determines which records or aggregates may cross the boundary and under what protections; once an authorized evidence stream reaches the operator, the next engineering problem is turning those records into continuous bias detection and production monitoring.

16.5.5 Bias detection and monitoring

Periodic pre-release audits cannot catch disparities that emerge between reviews when traffic or applicant populations shift. Suppose a nationally deployed credit model begins rejecting applicants from a specific ZIP code at twice its previous rate. Monitoring can expose that change before a scheduled audit, but the rate alone does not establish its cause. An investigator must check population changes, data quality, and decision policy against the approved fairness criteria. Bias detection turns formal definitions into operational telemetry, using slice-based analysis, which computes each metric separately for every slice of the population defined by attribute values, to identify subpopulations whose outcomes warrant examination.

16.5.5.1 Measurement costs and intervention points

The criteria of section 16.5.1 define what a disparity is and how to test one. Computing them continuously on a production system that serves thousands of predictions per second is where the engineering cost appears. That cost has computational, data, and organizational parts, and together they decide which metrics a system can afford to monitor.

Measurement overhead arises because computing group-specific metrics requires maintaining separate statistics for each protected group: for k groups and m metrics, this adds $\mathcal{O}(km)$ counters and $\mathcal{O}(km)$ statistical tests per evaluation cycle. Streaming counters can update cheaply per prediction while statistical tests run at a separate evaluation interval, so profiling the number of groups, windows, and tests, rather than request rate alone, determines whether sampling or asynchronous aggregation is needed. Those counters also depend on which fields the metric consumes, because error-rate criteria need outcome labels (Y) alongside sensitive attributes (G) across a sample that

supports the intended population claim, whereas demographic parity needs group membership and decisions without outcomes. Availability of those sensitive attributes requires an explicit governance and engineering choice among authorized direct collection, restricted audit records, encrypted or differentially private group statistics, or an offline sample. When inferred attributes are used instead, the audit must evaluate inference errors rather than treat the inferred groups as known. When the evidence cannot support a metric, the system must report that limitation rather than interpret missing measurements as fair behavior.

Selecting acceptable disparity thresholds then demands domain expertise and stakeholder input, since the four-fifths rule is a screening convention from U.S. employment selection rather than a universal legal test for lending or fraud detection. Thresholds must reflect context-specific harm assessments, statistical uncertainty, and documented legal rationale. Because distribution shift, feedback loops, or model updates can erode fairness over time, temporal stability further requires monitoring those thresholds across evaluation windows. Automated alerting (for example, "alert if $|\text{SPD}| > 0.15$ for 7 consecutive days") defines when an owner receives a signal. The persistence window reduces sensitivity to brief fluctuations at the cost of response delay, so its acceptability depends on the harm and exposure during that interval, as section 16.9.3 develops.

Offline and CI-stage audits are the first place these metrics run, and open-source toolkits such as Fairlearn package the per-group bookkeeping (Bird et al. 2020), as in the audit of listing 16.1. Applied to the running loan example, that audit's per-group table would report approval rates of 0.70 for Group A and 0.40 for Group B, the disparate impact ratio of 0.57 that falls below the 0.8 screening threshold in table 16.6. That comparison is an audit signal, not a lending-law verdict. The architecture must supply aligned decisions and group membership through an authorized audit path. Metrics that require outcome labels must use the relevant observed outcomes rather than substitute past decisions; an offline audit does not require exposing group attributes throughout runtime serving.

Listing 16.1: Bias Detection with Fairlearn: Illustrative audit excerpt for evaluating a loan-approval model across demographic groups. The excerpt assumes aligned arrays of observed outcomes, predictions, and demographic labels.

```
import numpy as np
from fairlearn.metrics import MetricFrame, selection_rate
from sklearn.metrics import accuracy_score, precision_score

def false_positive_rate(y_true, y_pred):
    y_true, y_pred = np.asarray(y_true), np.asarray(y_pred)
    negatives = y_true == 0
    if not negatives.any():
        return float("nan")  # No observed negatives: rate is unmeasurable.
    return ((y_pred == 1) & negatives).sum() / negatives.sum()

# Loan approval model evaluation across demographic groups
mf = MetricFrame(
    metrics={
        "approval_rate": selection_rate,
        "accuracy": accuracy_score,
        "precision": precision_score,
        "false_positive_rate": false_positive_rate,
    },
    y_true=loan_repaid_actual,  # observed repayment outcome, not past approval
    y_pred=loan_approvals_predicted,
    sensitive_features=applicant_demographics["ethnicity"],
)

# NaN means a group lacks negative examples, not that its FPR is zero.
# Display performance disparities across groups
print("Loan Approval Performance by Group:")
print(mf.by_group)
# Hypothetical output on the running loan example:
# approval_rate 0.70 for group A, 0.40 for group B
```

Where sensitive attributes are unavailable at inference time, as is common in mobile, federated, and embedded systems, group-specific thresholds may be infeasible. Restricted offline audits can still guide data changes, retraining, or a revised policy that does not require an attribute lookup on each request.

Fairness interventions may be applied at different points in the pipeline, but each comes with system-level implications. Preprocessing methods change the training distribution through sampling, reweighting, or augmentation. Their data requirements depend on the method; group-targeted reweighting needs group information, while feature augmentation also needs access to the affected examples. Versioning and lineage must preserve which transformation produced the training data. In-processing approaches instead embed fairness constraints directly into optimization. They require support for the chosen loss or constrained solver and may add training work or validation requirements.

↩ Feature versioning and the lineage records that trace each training input are established in section 12.4.4.

At the serving layer, postprocessing methods apply an approved decision rule to model outputs. Group-specific thresholds require the component choosing the threshold to obtain authorized group information; other policy rules may use different inputs. The dual-threshold policy in figure 16.5 ($\tau_{\text{thr},a} = 0.75$, $\tau_{\text{thr},b} = 0.45$) therefore needs a controlled attribute lookup and a record of the applied policy, but not propagation of the attribute through every routing and inference component. Access controls should limit that exposure. Validation must evaluate the resulting decisions, operational stability, and applicable restrictions on attribute use.

Multicalibration (section 16.5.2.3) can run in backend training infrastructure even when the resulting model serves on a constrained endpoint, but a training-time result does not establish calibration after deployment conditions change. If telemetry cannot support renewed subgroup audits, the operator must record that evidence limitation and restrict the approved operating envelope accordingly.

Mitigation also decays after launch. A fairness-aware model may degrade if retraining pipelines omit fairness checks, if logging cannot track subgroup outcomes, or if user feedback introduces biases absent from the training distribution, so retraining protocols need subgroup-labeled validation data under data-governance review. Deciding which disparity to remediate, and at what cost to other objectives, remains a stakeholder decision (section 16.9.4.3). The architecture's job is to make each disparity visible, in time, to the owner who must make that decision.

16.5.5.2 Real-time fairness monitoring architecture

Delivering that visibility without stalling live inference requires decoupling the serving path from the governance path (figure 16.6). A synchronous request path enforces access control, data minimization, and low-latency inference, while an asynchronous governance tier ingests bounded execution evidence and delayed labels to drive fairness auditing, post-hoc explanations, and owner-governed release gates.

The request path applies access controls and data minimization before inference and records only the evidence needed for the approved monitoring policy. Group-level fairness analysis and expensive explanations can run asynchronously; synchronous checks must fit the serving service-level objective (SLO). Delayed labels enter the analysis path when outcomes become available, and a named owner approves model or policy changes.

Real-time fairness monitoring updates prediction-only metrics such as demographic parity as decisions arrive, maintaining rolling statistics across protected groups. Label-dependent metrics such as equalized odds and equality of opportunity update when delayed outcomes or audited labels arrive. Groups whose labeled windows are too small must remain visible as having insufficient evidence, with a defined collection or escalation policy. How small is too small follows from the precision the monitor needs, because the uncertainty in a group's rate shrinks only with the square root of its labeled count; a group is under-supported when the confidence interval for its gap is wider than the tolerance the alert must resolve. The system flags supported disparities that exceed configured thresholds and separately reports gaps in coverage. Storage and compute requirements depend on the number of groups, decision segments, and retention windows.

The explanation engine generates evidence for the audience and decision under review. Its budget depends on the method and model; exact model-specific attribution may be feasible where exhaustive subset evaluation is not. Approximation, caching, selective generation, and asynchronous execution

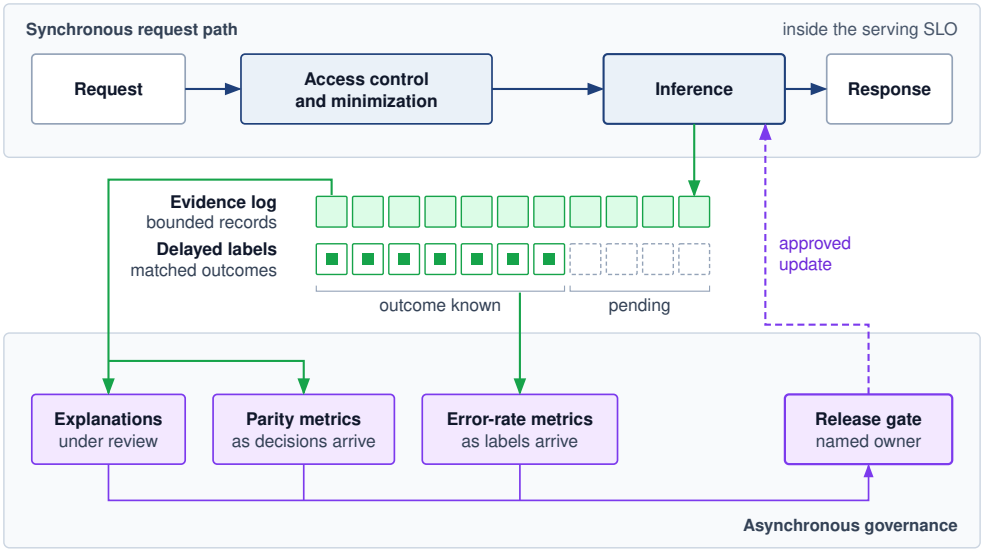

Figure 16.6: Production Responsible AI Architecture: The request path performs access controls, minimization, and inference. Bounded evidence collection feeds asynchronous fairness monitoring and explanation services. Delayed outcomes support label-dependent metrics; a governed update path returns approved changes to serving.

offer different cost and delivery-time trade-offs. The monitoring part of this architecture collects predictions and computes rolling group metrics with threshold alerts (listing 16.2). The listing omits delayed outcome ingestion, the full equalized-odds calculation, and explanation generation, so it is a partial implementation rather than a complete governance service.

This simplified pattern separates metric availability from metric values, because an unavailable comparison is not a zero disparity. A production monitor must also enforce required group coverage, sample-size policies, and freshness before interpreting a result. The overhead ranges summarized in table 16.5 explain why fairness monitoring, privacy transformations, and explanation generation need explicit capacity budgets. Request metadata collection may be synchronous, while metric aggregation, delayed-label analysis, and explanation generation can run asynchronously under their own freshness requirements. These overheads must be balanced against reliability and compliance requirements when designing production systems.

At the scale of a global production fleet, fairness monitoring becomes a substantial data engineering challenge. This sizing scenario assumes roughly 864 million inferences per day at 10,000 queries/s across 50 demographic subgroups, with 150 metrics (false positive rate, true positive rate, and calibration error for each of the 50 groups). A 1 percent sampling rate yields 8.64 million monitoring events daily. At an assumed 200 bytes per record, those events require approximately 1.7 GB per day before additional audit records or storage overhead. Actual record contents and retention follow the approved measurement and privacy requirements; outcome-dependent metrics also require labels that may arrive later. The monitoring infrastructure itself must therefore be reliable, secure, and cost-effective.

At a 5 percent false-alarm rate, 150 metrics raise about 7.5 spurious alerts per interval.

With 150 active metrics, an illustrative per-metric false-alarm probability of 5 percent would yield roughly 7.5 false alerts per evaluation interval if every tested null hypothesis is true. This expected count follows by summing false-alarm probabilities and does not require independence. Production thresholds must instead reflect subgroup denominators, delayed labels, repeated testing, confidence intervals, minimum detectable effects, and correction for multiple comparisons. A minimum detectable effect is the smallest gap that a group's sample size can reliably distinguish from noise, so it shows which groups are too small to monitor at the approved tolerance. A multiple-comparison correction, such as bounding the expected fraction of alerts that are false, tightens each

Listing 16.2: Production Fairness Monitoring: Simplified monitoring pseudocode for processing predictions and computing rolling fairness metrics. Delayed outcome ingestion is omitted.

```
import numpy as np

class RealTimeFairnessMonitor:
    # alert_threshold mirrors the |SPD| > 0.15 rule; persistence check omitted.
    def __init__(self, window_size=1000, alert_threshold=0.15, min_group=30):
        self.min_group = min_group  # Illustrative policy; set from precision needs.
        self.window_size = window_size
        self.alert_threshold = alert_threshold
        self.window = []

    async def process_prediction(self, prediction, demographics,
                                 actual_label=None):
        # Slide a fixed window over the stream, then re-check fairness.
        self.window.append((prediction, demographics, actual_label))
        self.window = self.window[-self.window_size :]
        metrics = self._compute_metrics()
        if (metrics["demographic_parity"] is not None
            and metrics["demographic_parity"] > self.alert_threshold) or (
            metrics["equalized_odds"] is not None
            and metrics["equalized_odds"] > self.alert_threshold
        ):
            await self._trigger_bias_alert(metrics)
        return metrics

    def _compute_metrics(self):
        # Bucket predictions and labels by protected group.
        by_group = {}
        for pred, demo, label in self.window:
            bucket = by_group.setdefault(
                demo.get("ethnicity", "unknown"), {"preds": [], "labels": []}
            )
            bucket["preds"].append(pred)
            if label is not None:
                bucket["labels"].append((pred, label))

        rates = [np.mean(d["preds"]) for g, d in by_group.items()
                 if g != "unknown" and len(d["preds"]) >= self.min_group]
        parity = max(rates) - min(rates) if len(rates) > 1 else None
        return {
            "demographic_parity": parity,
            # Requires sufficiently labeled outcomes per group; omitted here.
            "equalized_odds": self._equalized_odds_gap(by_group),
            "groups": {g: len(d["preds"]) for g, d in by_group.items()},
            "eligible_groups": len(rates),  # Compare with required group coverage.
        }

    def _equalized_odds_gap(self, by_group): ...
    async def _trigger_bias_alert(self, metrics): ...  # page and audit
```

test's threshold so that adding metrics does not multiply false alerts. A threshold crossing opens an investigation; it does not by itself prove discrimination or justify automatic retraining.

↩ False alarms, the spurious threshold crossings that accumulate as a fleet multiplies monitored metrics, are established in section 12.6.2.

In C^3 terms, this monitoring load is communication and coordination added to the fleet. Every sampled record must travel from a serving replica to the aggregation tier, and every alert must reach an owner who can act before the harm it signals compounds. The deployment tier sets how long that path is. A cloud service aggregates telemetry continuously, subject to storage and privacy budgets. An intermittently connected client keeps bounded local records until it next synchronizes, and a safety-critical edge system needs local checks that can trigger fallback within the decision deadline because the fleet cannot respond in time. A disconnected TinyML endpoint may deliver its evidence only at a servicing visit, and correcting its model may require a firmware update or

physical replacement. Those differences change the evidence delay and remediation path, not the requirement to name an owner for each monitored obligation.

Checkpoint 16.1: Putting bias detection into production

The bias-detection and monitoring sections turned static fairness definitions into live telemetry: slice-based detection with Fairlearn, CI-stage audits, production monitoring of demographic parity and equalized odds, the evidence paths that federated and on-device deployments require, and the distinction between observing a disparity and remediating it.

Detecting the disparity

- ☐ **Serving-path detection**: identify the request metadata needed to detect a rising rejection rate for one ZIP code, and explain how asynchronous aggregation can detect that change within an agreed monitoring delay.
- ☐ **Metric selection**: explain why the running loan example's approval rates of 70 percent for Group A and 40 percent for Group B do not select a fairness criterion. State the harm priority, label validity, and outcome evidence needed to choose one; then justify a criterion if reviewers prioritize avoiding denial to applicants who would repay.

From signal to action

- ☐ **Detection vs. remediation**: explain why a fairness alert with no remediation owner is evidence without action, and name what the system must add to close that loop.
- ☐ **Constrained deployments**: reason about why slice-based detection may be infeasible in a federated deployment where demographic labels cannot be centralized, and infer the implication for safeguard selection.

16.6 Privacy and Data Governance

Every fairness audit and monitor consumes demographic evidence, and that evidence is itself sensitive data. The threat models and privacy mechanisms of chapter 13 provide the technical controls; governance defines the evidence, ownership, and lifecycle interfaces that production systems require when inheriting them. A deletion workflow must identify which records and derived artifacts fall within an applicable obligation, such as an erasure requirement under the GDPR (European Parliament and Council of the European Union 2016). Removing a database row may leave replicas, backups, or trained models to address. Where removal of training influence is required, lineage, retraining, validation, deployment, and audit become part of the same workflow.

Governing these risks coordinates two classes of control that are easy to conflate. Privacy mechanisms bound specified information exposures, while machine unlearning (section 16.6.4) targets the influence a record retains in trained weights under a stated removal criterion. A privacy guarantee does not by itself remove a record from the training history, so the assurance case must distinguish these obligations and supply evidence for each. Robustness and generative-system controls receive the same evidence-and-ownership treatment in section 16.8.

16.6.1 Core data governance principles

That coordination starts at data ingestion, where a compliance rule demanding data minimization can collide directly with a fairness audit that needs a protected demographic field. Whereas technical privacy mechanisms constrain exposure and leakage, data governance determines whether collecting the field is justified, who may access it, and how long the audit may retain it. This decision connects privacy to fairness (section 16.5.1) and to transparency (section 16.7.1). Group-based monitoring needs evidence about group membership, but the audit design determines whether operators receive individual fields, restricted audit records, or aggregate statistics. Explainability methods can also reveal information about training examples, while detailed disclosure can conflict with individual privacy. Responsible AI systems must justify the evidence collected and limit its exposure while retaining enough information to evaluate the chosen requirements.

↩ Federated learning, encrypted computation, and synthetic data, each bounding who can observe raw records, are established in section 13.8.5.

Because training and personalization rely on personal data, privacy is a design constraint across the whole system lifecycle, and its central difficulty is the tension between data utility and individual protection. Rich, high-resolution datasets can enhance model accuracy and adaptability but also heighten the risk of exposing sensitive information, particularly when datasets are aggregated or linked with external sources. For example, large language models trained on broad text corpora have been shown to memorize specific strings that can later be retrieved through model queries or adversarial prompting (Carlini et al. 2021).

↩ The GPT-2 extraction attack, which recovered verbatim training records through ordinary generated outputs, is established in section 13.5.2.4.

Privacy risks also depend on how records are linked over time and across sources. A wearable device's heart-rate or location record can already be sensitive; combining physiological and behavioral signals can reveal additional details about a user's health, routines, or relationships. Limited visibility or control over processing, retention, and transmission makes it harder for users to understand or contest these uses.

Addressing these challenges requires understanding privacy as a system principle that entails robust data governance. This includes defining what data is collected, under what conditions, and with what degree of consent and transparency. Foundational data engineering practices, including data validation, schema management, versioning, and lineage tracking, provide the technical infrastructure for implementing these governance requirements. Responsible governance requires attention to labeling practices, access controls, logging infrastructure, and compliance with jurisdictional requirements. These mechanisms serve to constrain how data flows through a system and to document accountability for its use.

Weak data governance can obscure both quality defects and privacy exposures, but they require different evidence. In clinical applications, missing values or demographic skew can compromise predictive performance; neither alone establishes a privacy violation. Privacy assessment must also examine identifiability, access, and disclosure. Documentation should make these distinct risks traceable to the data practices that create them.

Privacy is not solely the concern of isolated algorithms or data processors; like fairness, it must be addressed as a structural property of the system. Decisions about consent collection, data retention, model design, and auditability all contribute to the privacy posture of a machine learning pipeline. This requires anticipating risks during inference and continuous operation as well as during training. Threats such as membership inference, which infers from a deployed model's outputs whether a record was in its training set, show why privacy safeguards must be embedded in both model architecture and interface behavior.

↩ Membership inference, its mechanism, and why an empirical audit gives only a lower bound on leakage are established in section 13.5.2.1.

Legal frameworks make privacy an enforceable design concern. The GDPR, the California Consumer Privacy Act (CCPA),[10] and Japan's Act on the Protection of Personal Information establish different obligations and individual rights within their respective scopes (European Parliament and Council of the European Union 2016; California Legislature 2023; Personal Information Protection Commission, Japan 2023). An implementation must identify which rules govern its collection, use, retention, and removal of data rather than treating consent or deletion as a uniform requirement across jurisdictions. That applicability decision belongs in the system's documented data-governance requirements.

Privacy mechanisms limit specified data exposures under their threat models; they do not establish that a system's decisions are safe. A model may satisfy its privacy requirements while still failing catastrophically on unexpected inputs, a complementary concern about *how systems fail* that section 16.8 takes up. Within privacy itself, the next decision is where the protected boundary sits in each deployment architecture.

[10] **CCPA (California Consumer Privacy Act)**: CCPA grants California residents rights over covered personal information, including a deletion right subject to statutory exceptions (California Legislature 2023). For ML serving infrastructure, deletion requests create an architectural constraint that must be planned from the data pipeline forward, because raw-record deletion does not automatically remove a record's influence from trained model parameters.

16.6.2 Privacy across deployment tiers

Training a model with differential privacy does nothing to protect a user whose raw prompt is copied into a plaintext serving log at the API gateway. Keeping raw records on a phone, by contrast, protects against central database breaches at the cost of hiding fleet-wide telemetry from operators. Privacy inherits the same centralization trade-off across deployment tiers (privacy row of table 16.4). Aggregating data in the cloud enables high-capacity modeling and monitoring but concentrates exposure to breaches and surveillance, so it must be paired with strong encryption, access control, and auditability. Keeping data local on mobile, federated, or TinyML clients reduces that central exposure but changes the evidence available to fleet operators. Privacy controls must cover local

storage, outputs, telemetry, and update channels; an endpoint with no update path needs controls that remain valid for its approved service life.

The constraint the table cannot show is that privacy claims must name their boundary. Differentially private training bounds training-data leakage through the released model; ordinary inference on fixed weights remains within that post-processing guarantee. It does not protect new private inputs copied into service logs or unrelated private-data accesses. Interface design, retention rules, rate limiting, and access control address those separate exposures. Opaque data collection can also violate user expectations, making consent surfaces and clear disclosure part of the privacy architecture.

Federated training draws three such boundaries (figure 16.7). Raw records stay on the client. Secure aggregation lets the server see only the sum of the client updates, and it hides individual updates only under its assumptions about collusion among parties and client dropout. Central differential privacy bounds what the released model reveals only if the curator that adds the noise is trusted. Each protection covers a different quantity against a different observer, so a claim about one says nothing about the others. The privacy unit, clipping rule, and protocol assumptions define each guarantee, and so does composition, the rule by which privacy loss adds up across training rounds (Bonawitz et al. 2017; Dwork and Roth 2014).

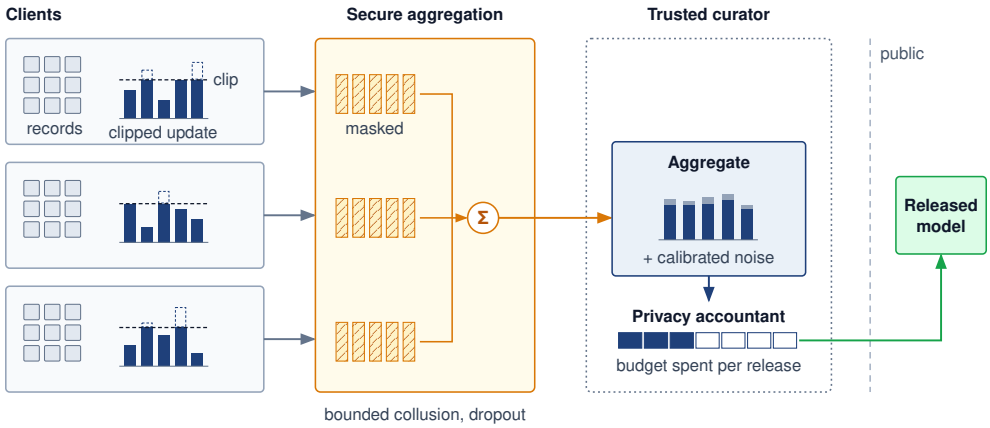

Figure 16.7: Privacy Boundaries in Federated Training: Each protection holds only at its own boundary. Raw records stay on clients, secure aggregation protects individual updates under its protocol assumptions, and central differential privacy bounds the released model relative to the specified privacy unit, provided the curator is trusted. Neither implies a universal ordering of privacy strength.

An architecture decides where the privacy boundary sits; the model decides how much crosses it. A service can keep raw records behind every access control and still release training data through the weights it serves, which is where the governance of privacy guarantees begins.

16.6.3 Governing privacy guarantees

A language model that regurgitates a patient's phone number from its training set or a classifier whose confidence scores reveal whether a specific customer was in the training cohort leaks private data through authorized outputs even when no database is breached. Governing a privacy guarantee starts from these output leaks, developed in section 13.5.2, and turns each into an audit obligation across data collection, model behavior, and user interaction. Because a membership-inference audit must state the attacker access it assumes and reads only as a lower bound on leakage, a formal bound has to come from the training mechanism. Differential privacy supplies such a bound. Its privacy budget ϵ limits how much the probability of any output can change when one privacy unit, the person or record the guarantee protects, is added or removed, except with a small probability δ. A privacy accountant totals the budget that training spends across all its steps.

For differential privacy, governance consumes the mechanism and accounting established in section 13.9. A differentially private stochastic gradient descent (DP-SGD) deployment must retain

the clipping, noise, sampling, and composition configuration that supports its reported bound (Abadi et al. 2016). Reviewers then evaluate that bound alongside utility, subgroup performance, and the private-data accesses outside training. Naming the method is insufficient evidence for approving the complete service.

Differential privacy introduces system-level trade-offs on the path where private data is accessed. Training noise can reduce utility, and attempts to recover it through additional data or optimization steps must be evaluated with the resulting privacy accounting and compute cost. If training occurs centrally, those costs do not automatically become extra work for an edge device performing ordinary inference on the released model. Local training or local randomization places different demands on the device. Data minimization, retention limits, and access controls remain separate requirements; feature obfuscation alone provides no differential privacy guarantee.

Systems Perspective 16.1: The price of privacy

Training a large language model on sensitive messages with DP-SGD reduces training-data extraction risk, but the protection is priced in compute and accuracy.

With the privacy unit, the neighboring relation (adding or removing one unit versus replacing one), and δ fixed, a smaller ϵ gives a stronger bound, and every change that lowers it (more noise, a smaller sampling rate, fewer steps) tends to cost accuracy or extra training. Section 13.9 derives these effects and lists the accountant inputs a reproducible DP-SGD result must record. Those recorded inputs, not the ϵ value alone, are what a governance review retains. The three configurations below are planning assumptions whose ϵ values an external accountant would have to confirm, as the CIFAR-10 accountant inputs of section 13.9.2.3 illustrate. Read them for the direction of each trade. The noise multiplier sets the noise scale as a multiple of the clipping norm. The stronger configuration doubles the noise multiplier of the weaker one but also trains for more epochs, and more steps by themselves raise ϵ; only the accountant can confirm that the extra noise more than offsets the extra steps.

Illustrative stronger-privacy configuration (assumed accountant result $\epsilon = 1$): The scenario assumes per-example gradients are clipped to norm $C = 1.0$, 40 percent of updates are affected by clipping, and the noise multiplier is $\sigma = 1.0$. It further assumes 3× more epochs, a training-cost increase from a GPT-3-scale baseline of $4.6M to approximately $13.8M, and an accuracy drop of 6 percent.

Illustrative weaker-privacy configuration (assumed accountant result $\epsilon = 8$): The scenario assumes a noise multiplier of $\sigma = 0.5$, training overhead of 30 percent, and an accuracy drop of 1 percent.

Illustrative weakest configuration (assumed accountant result $\epsilon = 10$): The scenario assumes less noise and less than 1 percent accuracy loss, with a weaker formal privacy bound than the other two configurations under otherwise comparable accounting assumptions.

These scenarios compare formal privacy bounds, utility, and training cost; none of those quantities alone establishes user trust. Once a fixed model has been trained with differential privacy, ordinary inference is post-processing and does not consume additional training-data privacy budget. Additional privacy loss must be accounted for when a postdeployment mechanism makes new accesses to private data or releases additional differentially private statistics. The key engineering decision is to document and account for every private-data access and communicate the resulting trade-offs to users and regulators.

Released datasets need a boundary of their own. A k-anonymous release (section 13.8.5.3) hides each record among at least k_{anon} that share its quasi-identifiers, a property of a specified release rather than of one transformed request. Differential privacy likewise requires a defined release mechanism, privacy unit, neighboring relation, and privacy accountant; noise injection alone does not establish a guarantee.

Privacy enforcement also depends on infrastructure beyond the model itself. Data collection interfaces must support informed consent and transparency. Logging systems must avoid retaining sensitive inputs unless strictly necessary, and must support access controls, expiration policies, and

auditability. Model serving infrastructure must be designed to prevent overexposure of outputs that could leak internal model behavior or allow reconstruction of private data. These system-level mechanisms require close coordination between ML engineering, platform security, and organizational governance.

Privacy must be enforced throughout the machine learning lifecycle. Retraining pipelines must respect the approved deletion and data-use policies. Monitoring should collect only the identifying information justified by its purpose, with access and retention controls appropriate to that information. Privacy-aware telemetry collection and protected per-user audit trails can support investigation and remedy without making sensitive records broadly available in logs or dashboards.

The privacy assurance case therefore spans the whole pipeline. The formal privacy accountant covers the private-data accesses in its stated mechanism; interface controls, retention policies, and access records support separate claims about the surrounding service. Keeping these claims distinct makes gaps visible when a model, telemetry path, or data source changes.

Privacy preservation techniques create complex sociotechnical tensions that extend well beyond technical implementation. Differential privacy mechanisms may reduce model accuracy in ways that disproportionately affect underrepresented groups, creating conflicts between privacy and fairness objectives. These challenges require ongoing stakeholder engagement as detailed in section 16.9.4.3, where organizations must navigate competing values around data control, personalization, and regulatory compliance. Even a strong training-time privacy bound, however, treats the training set as fixed, leaving open what happens when a user revokes consent or demands erasure after the model is already in production.

16.6.4 Machine unlearning

Deleting a user's row from a production database after an erasure request under GDPR or CCPA (European Parliament and Council of the European Union 2016; California Legislature 2023) leaves their influence intact inside a 175-billion-parameter model trained on last month's snapshot. Because privacy preservation does not end at training time, machine unlearning supplies a technical response when the approved deletion policy also requires addressing retained model state. The legal requirement and the evidence of removal must be specified separately, since retained influence alone does not establish that every deployment violates a deletion obligation. When the approved response does require removing training influence, the architecture must support the chosen method and its verification: an endpoint without sufficient data, training state, or compute needs a replacement model trained elsewhere, whereas a device with the required resources may perform the work locally. Placement follows those resource and state requirements rather than the deployment label alone.

Three update strategies differ in what they recompute and what removal they guarantee (figure 16.8). Retraining from scratch on the retained data, the training set with the forget set D_{forget} removed, supplies a reference for exact removal relative to that procedure. Gradient ascent on the forget set increases its loss through updates of the form $\theta + \eta\nabla\mathcal{L}$, but increased loss alone does not demonstrate removal. Sharded, isolated, sliced, and aggregated (SISA) training partitions data into isolated shards S_1, S_2, S_3 whose constituent models are combined by aggregating their predictions at serving time. Within each shard, the data is further ordered into sequential slices introduced incrementally during training, with a checkpoint saved before each slice is added. A deletion therefore affects only the constituent model trained on the record's shard, and rewinding that shard to the checkpoint saved just before the affected slice avoids replaying earlier slices.

Full retraining is a useful reference, but its feasibility depends on access to the permitted remaining data, the training procedure, and sufficient compute within the response deadline. A retention policy may remove data that retraining would need; a large training run may exceed the available budget. An alternative unlearning method must state its own data and state requirements rather than assume that avoiding full retraining also eliminates those dependencies.

Machine unlearning seeks to remove specified training influence without repeating the entire training run. Cao and Yang (2015) proposed converting suitable learning algorithms into a summation form, in which the trained model depends on the data only through sums over training examples, allowing a deletion to subtract its terms from those sums and recompute the model. SISA instead localizes retraining to affected shards (Bourtoule et al. 2021). It still needs the retained

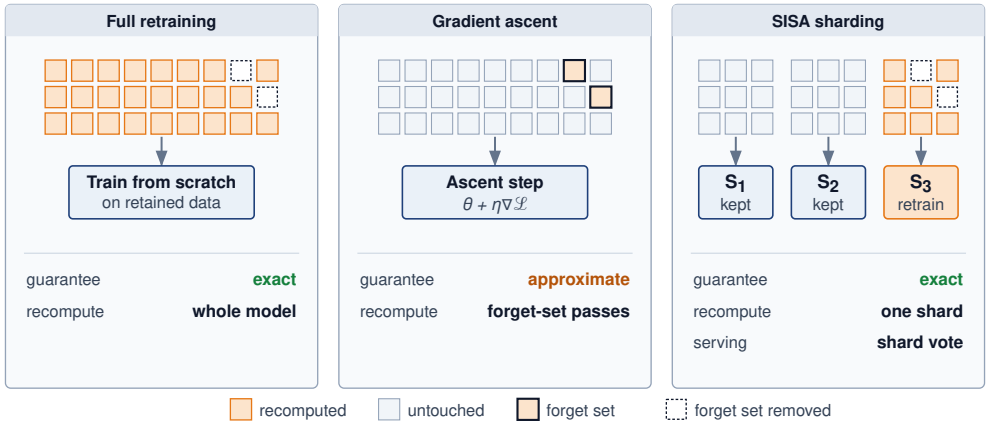

Figure 16.8: Model Update Strategies: Three approaches to removing a data point's influence. Full retraining gives exact unlearning at the highest cost; gradient ascent provides an approximate update by maximizing loss on the forget set; and SISA training partitions data into shards so only the affected shard is retrained. The accompanying cost model quantifies the full-retraining and SISA cases.

data and training state required by that procedure. The removal claim must specify its reference and verification method; a loss increase on deleted examples is different evidence from retraining without them. Method selection therefore depends on retained state, model quality, operational cost, and the removal guarantee the deployment must support.

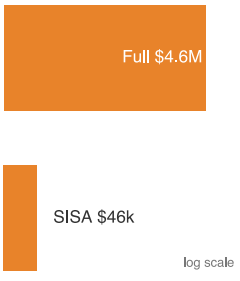

For one isolated erasure, SISA cuts retraining cost about 100 times; a daily batch that touches most shards erodes the saving.

Napkin Math 16.6: The cost of forgetting

A user requests erasure of personal data under GDPR Article 17 ("Right to Erasure"). This scenario assumes that the approved response includes removing its influence from a model trained on 1 TB of data; that model-removal requirement is a scenario premise, not a conclusion derived from invoking the article.

Baseline (full retraining): In this sizing scenario, a 175B-parameter model at GPT-3 scale retrains on 1,024 NVIDIA A100 (80 GB SXM4) GPUs (dense FP16 $R_{\text{peak}} = 312$ TFLOP/s per accelerator at an assumed model FLOPs utilization (MFU) of about 32 percent) for approximately 36 days at a cost of roughly $4.6M.

Engineering fix: SISA training partitions data into $K = 100$ independent shards, training 100 sub-models. For one datum confined to one shard, the example retrains that shard (1 percent of data). Assuming cost and elapsed time scale linearly with the data fraction at the same resource allocation gives $46,000 and approximately 8.6 hours. This simplified estimate omits slicing, checkpoint storage, validation, and deployment overhead; it is not a measured GPT-scale SISA result.

Trade-off: The scenario assumes an accuracy drop of 3–7 percent; the cost calculation does not establish that quality result. Serving must evaluate and aggregate 100 sub-models, with latency depending on parallelism and aggregation. Batch removal cost depends on how many of the 100 shards are touched. If the 1,000 requests/day arrive independently and uniformly across shards, a daily batch is likely to touch nearly all of them, eroding the sharding-only saving. Concentrated requests, batching policy, and saved slice states change the work. The operating comparison must therefore include the request distribution and response deadline, rather than the shard count alone.

Those cost and batching trade-offs explain why statutory erasure rights under GDPR, CCPA, and similar laws (European Parliament and Council of the European Union 2016; California Legislature 2023) require context-dependent engineering triage rather than a universal command to retrain

every model from scratch. Where model state itself retains personal data, as highlighted by incidents in which generative models reproduced personal content or copyrighted text, machine unlearning provides a technical strategy for reducing or removing a deleted record's influence within a feasible compute budget.

From a systems perspective, machine unlearning introduces nontrivial architectural and operational requirements. Systems must be able to track data lineage, including which datapoints contributed to a given model version. This often requires structured metadata capture and training pipeline instrumentation. Additionally, systems must support user-facing deletion workflows, including authentication, submission, and feedback on deletion status. Verification may require maintaining versioned model registries, along with mechanisms for testing the stated removal guarantee against its reference training procedure. Such tests can check that membership inference no longer distinguishes the deleted records from records the model never saw, and that the updated model's outputs match those of a model retrained without the deleted records within a stated tolerance. These operations must span data storage, training orchestration, model deployment, and auditing infrastructure, and they must be robust to failure or rollback.

Resource-constrained deployments amplify these challenges further. A TinyML device may lack the compute needed for retraining while still accepting a replacement model trained elsewhere. Unlearning then requires fleet inventory, authenticated updates, and evidence that affected versions have been replaced or retired. If no update path exists, operators must plan for physical servicing or withdrawal; initial data minimization reduces exposure but does not erase influence already present in deployed weights. Even in cloud-based systems, unlearning must coordinate model replicas, snapshots, and logs.

The need to remove training influence should be assessed during architecture and policy design, when teams can still plan the required lineage, retained state, and update paths. Retrofitting those capabilities after deployment may require reconstructing provenance or rebuilding training infrastructure. The assurance case should specify which removal obligations apply, how the chosen method will be evaluated, and how approved replacement models reach deployed instances.

Machine unlearning represents a shift in privacy thinking, from protecting what data is collected to controlling how long that data continues to affect system behavior. Exercise of any data right or challenge to an automated outcome, however, depends on whether users, clinicians, and auditors can see how the system works and why it produced a specific decision.

16.7 Explainability and Transparency

A rejected loan applicant, a clinician questioning a risk score, and an auditor reconstructing a contested decision each need a different account of the same model output, on a different deadline. Serving them requires disclosure about the system as a whole, explanation methods whose cost and fidelity fit the deployment, and an appeal path through which the affected person can act on what is disclosed.

16.7.1 The need for transparency

When ProPublica audited the COMPAS recidivism-risk score used in U.S. courts, it found that Black defendants were falsely flagged as future offenders at nearly twice the rate of white defendants (Angwin et al. 2016).[11] Meanwhile, the vendor's trade-secret claims prevented defendants, judges, and external investigators from inspecting the scoring formula or training provenance. Machine learning models often operate as opaque "black boxes," producing outputs that users, developers, and regulators struggle to scrutinize. In high-stakes domains such as criminal justice, healthcare, and finance, that opacity blocks investigation and remediation even when external outcome audits expose a disparity.

The COMPAS case separates the two properties defined among the core principles. What investigators lacked was transparency, meaning disclosure of data sources, training procedures, evaluation protocols, and known limitations. Explanations of individual scores would have been a different kind of evidence, about particular decisions rather than the system. An attribution describes how a method assigns influence under its assumptions rather than supplying a causal account, a normative justification, or a path to recourse. Section 16.7.2 develops the methods, their costs, and their limits; the principle established here is that both properties are obligations the architecture must provision.

[11] **COMPAS (Correctional Offender Management Profiling for Alternative Sanctions)**: ProPublica's 2016 analysis found Black defendants were falsely flagged as future criminals at nearly twice the rate of white defendants (45 percent vs. 23 percent false positive rate) (Angwin et al. 2016). The proprietary, black-box nature of the system limited full inspection of model internals and training data, making the source of the observed error-rate disparity harder to investigate. Error-rate parity and predictive parity are distinct criteria, so the evaluation must name the criterion rather than treating one rate as a complete fairness verdict.

The importance of explainability and transparency extends beyond technical considerations to legal requirements. In many jurisdictions, these principles are legal obligations rather than optional guidance. Under the General Data Protection Regulation (GDPR), solely automated decisions with legal or similarly significant effects trigger Article 22 safeguards and related transparency duties, including meaningful information about the logic involved under access and notice provisions, subject to statutory exceptions (European Parliament and Council of the European Union 2016; European Data Protection Board 2018).[12] Similar regulatory pressures in other domains reinforce the need to treat explainability and transparency as core architectural requirements.

Implementing these principles requires anticipating the needs of different stakeholders, whose competing values and priorities are examined comprehensively in section 16.9.4.3. Designing for explainability and transparency therefore requires decisions about how and where to surface relevant information across the system lifecycle.

Transparency and explainability also support system reliability over time. As models are retrained or updated, mechanisms for interpretability and traceability allow detection of unexpected behavior, enable root cause analysis, and support governance. Embedded into the structure and operation of a system, these mechanisms provide the foundation for trust, oversight, and alignment with institutional and societal expectations.

While transparency and explainability enable stakeholders to understand system behavior, they do not guarantee that this behavior is equitable. A model can be fully transparent about how it makes decisions while still systematically disadvantaging certain groups, which is why the fairness criteria of section 16.5.1 remain a separate test. Fulfilling the explainability obligation itself requires choosing between models whose decision structure can be inspected directly and post-hoc methods that estimate feature influence after training.

12 | **GDPR (General Data Protection Regulation) Article 22**: Article 22 gives data subjects protections against decisions based solely on automated processing that produce legal or similarly significant effects, with exceptions for contract necessity, authorization by law, or explicit consent plus safeguards (European Parliament and Council of the European Union 2016). Related notice and access provisions, including Article 15, require meaningful information about the logic involved for qualifying automated decision-making, and EDPB guidance treats those duties as part of a broader profiling and automated-decision governance regime (European Data Protection Board 2018).

16.7.2 Models of interpretability

A loan officer using a traditional rules-based system can identify the rule that produced a rejection: "The debt-to-income ratio exceeds 40 percent." A neural network, however, produces a score through many interacting operations for which no single rule supplies the decision rationale. Explainability and interpretability provide bounded evidence about model behavior, such as feature influence, sensitivity, representative examples, or counterfactual changes. That evidence can support investigation and communication, but it becomes a justification only when the institution connects it to a valid decision policy, and it becomes recourse only when the user can take a feasible action that the system will honor.

Explanation evidence plays a role in system validation, error analysis, incident investigation, user communication, and some compliance workflows. In high-stakes domains such as healthcare, financial services, and autonomous decision systems, explanation methods can generate hypotheses about spurious correlations that engineers must test independently. For instance, a saliency[13] or sensitivity analysis might indicate that a diagnostic model responds to image artifacts rather than medical features; controlled evaluation must then determine whether that dependence is causal and operationally significant. For qualifying automated decisions under GDPR-style access and notice provisions, systems may need to provide meaningful information about the logic involved, reinforcing the need for systematic support for explanation, documentation, and contestability.

Two scopes of explanation place different demands on the system. Local explanations characterize influences or sensitivities around an individual prediction, so their cost depends on which decisions are explained and whether results can be reused. Global explanations characterize broader model behavior, using model structure, representative inputs, or aggregated decision evidence. They may be computed offline or updated as new evidence arrives. Scope describes what the explanation covers; the method and delivery requirement determine where and when it runs.

Explainability methods can be broadly categorized based on when they operate and how they relate to model structure. Post-hoc methods are applied after training, but they are not necessarily black-box methods. Model-agnostic approaches such as LIME (Ribeiro et al. 2016) and KernelSHAP can operate through model queries, whereas input gradients, integrated gradients[14] (Sundararajan et al. 2017), and gradient-weighted class activation mapping (GradCAM)[15] (Selvaraju et al. 2017) require differentiable-model access to gradients or internal activations. Other SHAP implementations (Lundberg and Lee 2017) may exploit model structure.

13 | **Saliency maps**: Gradient-based explanation that highlights which input regions most influenced a prediction by computing a backward pass, the same infrastructure used for training. Saliency methods are often cheaper than perturbation-heavy or subset-enumeration methods, making them attractive for edge and mobile deployments when the model supports gradients. The trade-off is that raw gradients can be noisy and may highlight input artifacts rather than meaningful features, so smoothing or repeated-gradient variants trade additional compute for more stable visualizations.

Post-hoc approaches are widely used in image and tabular domains, where explanations can be rendered as saliency maps or feature rankings. To illustrate how SHAP attribution works in practice, consider a trained random forest model predicting loan approval (approve $= 1$, deny $= 0$) based on three features: `income`, `debt_ratio`, and `credit_score`. For a specific applicant who was denied, with income of $45,000, debt ratio of 0.55 (55 percent of income goes to debt), and credit score of 620, the model outputs $\Pr(\text{approve}) = 0.08$, that is, denial with probability 0.92. SHAP values, based on Shapley values from cooperative game theory,[16] measure each feature's contribution to moving the approval probability from a baseline (the expected prediction over a declared background distribution, here $\Pr(\text{approve}) = 0.50$) to this individual prediction.

The contribution of a feature depends on which other features are already known, so adding income first and adding it last can shift the prediction by different amounts. SHAP resolves that ambiguity by averaging a feature's marginal contribution over every order in which the features could be added, under a specified value function. The value function states what the model outputs when some features are unknown; here it averages the prediction over the declared background distribution, which sets the 0.50 baseline. In this illustrative decomposition the averaged contributions are -0.05 for `income`, -0.25 for `debt_ratio`, and -0.12 for `credit_score`.

Shapley values are additive, so the contributions sum from the baseline to the prediction: $0.50 - 0.05 - 0.25 - 0.12 = 0.08 = \Pr(\text{approve})$, and therefore $\Pr(\text{deny}) = 0.92$. Under this attribution setup, debt ratio has the largest assigned contribution to the output, followed by credit score and income. This ranking is a hypothesis about model sensitivity under the chosen background distribution, not evidence that debt ratio caused the decision or that lowering it below 40 percent would flip the output. A recourse claim requires a separately validated counterfactual that respects feature dependencies, immutable attributes, feasible user actions, and the decision policy that will apply when the user returns.

The computational cost depends on the attribution algorithm. Enumerating all subsets would evaluate $2^3 = 8$ coalitions in this example and $2^{20} \approx 1$ million for twenty features. A random forest need not use enumeration: TreeSHAP exploits tree structure to compute exact attributions under its supported value-function assumptions. KernelSHAP estimates contributions from sampled coalitions, while DeepSHAP propagates approximations through network components. Gradient-based attribution can need fewer model evaluations still, but only where the deployed runtime supports gradient computation, and lower cost does not establish explanation fidelity. Explanation capacity must therefore be budgeted for the chosen method rather than inferred from feature count alone.

Another post-hoc approach involves counterfactual explanations, which estimate how a model output changes under specified input modifications. These are relevant for decision-facing applications such as credit or hiring systems, but a mathematical output flip is not automatically recourse. A useful counterfactual must distinguish mutable from immutable features, preserve domain constraints and feature dependencies, identify an action within the user's control, and remain valid under the institution's future decision policy (Wachter et al. 2017). Counterfactual generation and validation therefore require domain models, policy versioning, and operational commitments in addition to constrained optimization.

A third class of techniques relies on concept-based explanations, which attempt to align learned model features with human-interpretable concepts. For example, testing with concept activation vectors (TCAV) represents user-defined concepts as vectors in activation space and tests how sensitive a model's predictions are to those concepts in its internal representation (B. Kim et al. 2018). These methods are especially useful in domains where subject matter experts expect explanations in familiar semantic terms. However, they require training data with concept annotations or auxiliary models for concept detection, which introduces additional infrastructure dependencies.

Post-hoc methods answer method-specific questions about model behavior. Even an exact attribution under a chosen value function does not recover a causal rationale or justify the decision policy. Approximate methods add estimation or surrogate error to those interpretive limits. Evaluation must therefore establish both what the method computes and whether that evidence serves the intended debugging, communication, or review task.

Inherently interpretable models expose a decision structure that the intended audience can inspect. Small decision trees, short rule lists, and sparse linear models with meaningful features can support

[14] **Integrated gradients**: An attribution method that integrates gradients along the straight path from a baseline input to the actual input. Sensitivity requires a nonzero attribution when a single changed feature accounts for a changed prediction; implementation invariance requires identical attributions for functionally equivalent models (Sundararajan et al. 2017). The authors argued that most earlier attribution methods violate at least one of these axioms. These properties concern the exact method, while numerical integration introduces an approximation whose accuracy depends on the sampling of the path. The deployment must therefore budget repeated gradient evaluations and check approximation error rather than treating a fixed step count as sufficient for every model and input.

[15] **GradCAM (gradient-weighted class activation mapping)**: Selvaraju et al. (2017) generalized Class Activation Mapping to any convolutional neural network (CNN) architecture by using gradients flowing into the final convolutional layer to produce spatial importance maps. Because it reuses gradient computation, Grad-CAM can be far cheaper than perturbation-heavy explanation methods, but whether it fits real-time medical imaging or autonomous-vehicle pipelines depends on the model, hardware, batch size, and explanation-frequency budget.

[16] **Shapley values**: From cooperative game theory, Shapley values fairly distribute a payoff among players based on marginal contributions across all possible orderings. In ML explainability, features are "players" and the prediction is the "payoff." The mathematical guarantees (efficiency, symmetry, null player) make SHAP a widely used attribution framework with clear axiomatic appeal, but exact subset enumeration grows as $2^{n_{\text{feat}}}$ for n_{feat} features, which is why production systems often rely on model-specific algorithms, sampling approximations, caching, or asynchronous explanation paths (Lundberg and Lee 2017).

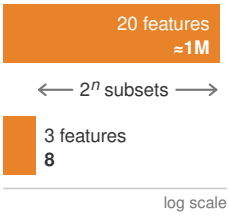

Exact subset enumeration grows exponentially; model-specific algorithms can avoid it.

this goal; a large tree or high-dimensional linear model may not. Nearest-neighbor comparisons likewise require an understandable similarity metric and representative examples. Interpretability therefore depends on structure, features, and audience, not merely the model-family label. In high-stakes applications, model selection should evaluate whether a sufficiently accurate, inspectable model meets the task before assuming that post-hoc explanations are necessary (Rudin 2019).

Six common model families fall along an interpretability continuum in three groups (figure 16.9). Sparse linear models (explained via coefficients) and small decision trees (via split paths) are interpretable by design. Prototype and concept models are hybrid architectures explained via exemplars. Tree ensembles (via TreeSHAP), deep networks (via gradients and SHAP), and foundation models (via probing, which trains a small classifier to read a property from internal activations) require post-hoc explanation. While this ordering provides a structural heuristic across model families, dimensionality and scale can shift boundaries; a million-node decision tree or dense linear regression can prove just as opaque as a deep neural network. This distinction is central to choosing an appropriate model for a given application, particularly in settings where regulatory scrutiny or stakeholder trust is essential.

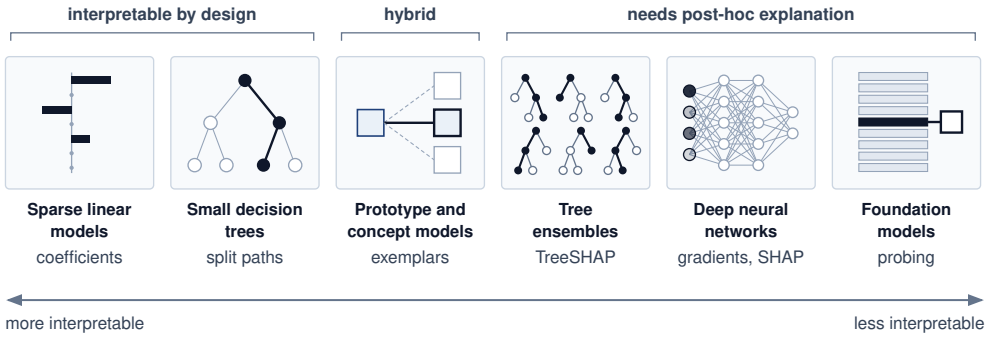

Figure 16.9: Model Interpretability Spectrum: A heuristic ordering by model family. Small linear models and shallow decision trees can offer transparent reasoning, while neural networks typically require post-hoc explanation techniques; size, feature meaning, and audience can move any model along the spectrum. This distinction guides model selection based on application needs, prioritizing transparency in regulated domains or when stakeholder trust is important. The label under each family names the handle an explanation typically uses.

Hybrid approaches aim to combine the representational capacity of deep models with the transparency of interpretable components. Concept bottleneck models (Koh et al. 2020), for example, first predict intermediate, interpretable variables and then use a simple classifier to produce the final prediction. ProtoPNet models (Chen et al. 2019) classify examples by comparing them to learned prototypes, offering visual analogies for users to understand predictions. These hybrid methods are attractive in domains that demand partial transparency, but they introduce new system design considerations, such as the need to store and index learned prototypes and surface them at inference time.

Work such as Olah et al. (2020) and Geiger et al. (2021) applies mechanistic interpretability to reverse-engineer the internal operations of neural networks. This line of work, inspired by program analysis and neuroscience, attempts to map neurons, layers, or activation patterns to specific computational functions. Claims about such mechanisms require instrumentation and tests that support the proposed interpretation. These research methods can contribute evidence about internal computation, but do not establish a complete audit of every deployed model.

From a systems perspective, explainability introduces a number of architectural dependencies. Explanations must be generated, stored, surfaced, and evaluated within system constraints. The required infrastructure may include explanation APIs, memory for storing attribution maps, visualization libraries, and logging mechanisms that capture intermediate model behavior. Models must often be instrumented with hooks or configured to support repeated evaluations, particularly for explanation methods that require sampling, perturbation, or backpropagation.

The sociotechnical challenges of explainability center on the gap between technical explanations and human understanding. While algorithms can generate feature attributions and gradient maps, stakeholders often need explanations that align with their mental models, domain expertise, and decision-making processes. A radiologist reviewing an AI-generated diagnosis needs explanations that reference medical concepts and visual patterns, not abstract neural network activations. This translation challenge requires ongoing collaboration between technical teams and domain experts to develop explanation formats that are both technically accurate and practically meaningful. Explanations can shape human decision-making in unexpected ways, creating new responsibilities for how explanatory information is presented and interpreted. That gap between a technically valid explanation and a usable one is the first instance of a broader problem. Every control in this chapter operates inside human workflows and institutions that no model-level method governs.

16.7.3 Explanation placement

Bridging that gap under production deadlines starts with where the explanation computation can afford to run. Running sampled KernelSHAP inline on a microcontroller or a 10 ms control loop will miss the actuation deadline before a single attribution finishes, whereas deferring every attribution to an offline batch queue leaves a rejected applicant with no immediate reason for the outcome. Computational capacity, latency budget, interface design, and data access constrain both the explanation method and its placement. A cloud worker can execute expensive SHAP or LIME attribution; a constrained endpoint may use a cheaper gradient-based saliency map if its runtime supports gradients. Devices unable to generate explanations locally may preserve bounded decision records for later analysis by a backend or during servicing, the per-tier split that the explainability row of table 16.4 records.

Audience and timing further constrain explanation design. An end user may need a concise summary ("elevated heart rate during sleep"), while a developer or auditor may need attribution maps, concept activations, or decision traces. The required evidence depends on what each person must understand or do. A drone or industrial controller must keep explanation work within its control deadline or move that work off the critical path. Financial and clinical workflows likewise need a specified delivery deadline, which may allow analysis after the automated decision but before review or appeal. Planning for the audience and deadline, within the deployment's resource limits, makes interpretability a design constraint.

Storage requirements also scale significantly with explanation needs. Storing SHAP values for tabular data requires approximately 4–8 bytes per feature per prediction, so monthly storage is 30 million predictions times the feature count times 4–8 bytes for a system logging 1 million predictions daily. Gradient attribution maps for images can require 1–10 MB per explanation depending on resolution; logging every image explanation at the same 1 million-per-day rate would require roughly 30–300 TB per month unless explanations are sampled, compressed, or retained for a shorter window. These volumes require careful data lifecycle management and retention policies.

Compression and optimization techniques also affect explainability. Pruning, quantization, and architectural simplifications often used in TinyML or mobile settings can distort internal representations or disable gradient flow, degrading the reliability of attribution-based explanations. In such cases, interpretability must be validated postoptimization to ensure that it remains meaningful and trustworthy. Even an explanation that survives optimization and fits its latency and storage budget, however, protects the user only if the surrounding service provides a way to act on it.

16.7.4 Transparency and contestability

Showing a rejected loan applicant a feature-attribution chart accomplishes nothing if the app offers no button to dispute an error in the credit file, no human reviewer with authority to overturn the rejection, and no immutable snapshot of the inputs that produced the score. Explanations serve accountability only when stakeholders can both understand and challenge system decisions. Transparency allows users, developers, auditors, and regulators to inspect how a system functions, assess its limitations, and identify sources of harm, but disclosure alone is not enough. In high-stakes domains, individuals and institutions need contestability: the operational ability to interrogate, appeal, and reverse an automated decision.

Documentation tools such as model cards, standardized reports published with a trained model, and datasheets for datasets, their counterpart for a dataset, formalize system-level disclosure by recording how models are trained, what data they rely on, what assumptions are embedded in their design, and what known limitations affect their use. Those artifacts improve governance, support compliance, and inform user expectations, yet static documentation still leaves affected individuals without interface tools, procedural access, or institutional support when a specific decision harms them.

To move from transparency to contestability, machine learning systems must integrate mechanisms for explanation, recourse, and feedback. Explanation provides understandable reasons for system outputs, tailored to the operational context and needs of the person receiving them. Recourse identifies feasible changes through which individuals could receive a different outcome under the stated decision policy. A proposed change needs to be actionable for that person; a counterfactual prediction alone does not guarantee an outcome after the model or policy changes. Feedback enables users and domain experts to dispute outcomes, report errors, or signal concerns, routing those signals into oversight processes and model updates.

These mechanisms can be missing from the user-facing workflow even when internal documentation exists. A mobile loan application might display a rejection without explaining how to submit additional information or reach an appeal reviewer. An external appeal process can still exist, but users need a discoverable route to it and access to the evidence relevant to their case. Similarly, a clinical model may produce a risk score without giving the physician enough context to question its use for a particular patient. If subgroup errors are neither observable nor contestable, the resulting harm becomes harder to diagnose and correct.

From a systems perspective, enabling contestability requires coordination across technical and institutional components. Models must expose sufficient information to support explanation. Interfaces must surface that information in a usable and timely way, so that organizational processes can review feedback, respond to appeals, and update system behavior. Logging and auditing infrastructure must also track user interventions and override decisions alongside model outputs. Human overrides and decision abstention can support this process only when qualified reviewers have capacity, evidence, and authority to resolve the deferred cases within the required deadline.

Implementing contestability imposes concrete infrastructure costs that scale with system throughput and complexity. Storing the necessary metadata to reconstruct a decision (input features, model version, and decision thresholds) requires persistent storage whose footprint depends on feature dimensionality, explanation payloads, and retention windows. Generating on-demand explanations using Shapley values or counterfactuals can add enough latency that contested decisions often need asynchronous processing queues to preserve serving service level agreements (SLAs). Meeting record-keeping and automatic-logging obligations for high-risk systems under frameworks such as the EU AI Act requires storage, provenance, retention controls, and oversight capacity to be budgeted as part of the inference fleet rather than as a separate policy artifact (European Parliament and Council of the European Union 2024).

War Story 16.1: Apple Card: A systems reconstruction of contested decisions (2019)

Context: In 2019, Apple and Goldman Sachs faced intense public scrutiny after a customer reported a credit limit 20× that offered to his wife despite joint tax filings. Other customers raised similar concerns about unequal terms (New York State Department of Financial Services 2021).

Evidence boundary: The public record documents complaints, the investigation, the regulator's finding, and changes to appeal policy. It does not disclose the underwriting pipeline's latency target or internal explanation architecture. The systems account below is therefore an explicit reconstruction, not a report of Apple or Goldman Sachs implementation details.

Reconstructed mechanism: A low-latency underwriting service that returns only a limit and generic adverse-action reason leaves support staff unable to inspect the decision snapshot, distinguish data errors from policy effects, or route a contested outcome to a reviewer. The missing interface is broader than feature attribution. It includes versioned input provenance,

policy and model versions, reasons appropriate to the decision process, appeal state, and authority to correct an error.

Observed impact: The New York Department of Financial Services found no violation of fair-lending law in its investigation, but the complaints still exposed a trust and contestability failure. Customers and frontline staff lacked a satisfactory account of how apparently inconsistent outcomes could be reviewed (New York State Department of Financial Services 2021).

Systems lesson: High-stakes decision systems need a contestability interface, not merely an attribution API. A defensible architecture preserves the decision-time evidence needed for audit, generates appropriate reasons without claiming that attribution is causation, routes appeals to an authorized human, records remediation, and tests the full workflow under its latency and retention constraints.

The Apple Card reconstruction names the pieces such a system lacks. One architectural pattern that supplies them is a contestability stack, analogous to distributed tracing in microservices, that coordinates four coupled stages across the operational pipeline. Decision provenance preserves a verifiable link from each output back to its governing policy, model version, and input snapshot, protected by authenticated logs or cryptographic integrity checks. Explanation generation budgets explanation methods tailored for appeals, audits, and routine communication, offloading compute-intensive methods to asynchronous workers while keeping lightweight techniques on the request path. When decisions are challenged, appeal routing directs contested cases to qualified human reviewers equipped with domain context. Closing the loop, outcome tracking records whether the appeal upheld or overturned the machine decision, feeding empirical discrepancies back into evaluation and retraining. Missing decision-time records can prevent investigators from reconstructing a contested decision, even if aggregate debugging remains possible, so the records to retain follow from what an investigation and a remedy will need.

The stack need not run where the decision runs. An edge or TinyML endpoint can hand its decision records to a backend review process at the next synchronization or servicing visit, which lengthens the appeal's evidence path without removing the obligation to offer one. In every tier, transparency supports accountability only when users and institutions can use the disclosed information to question and correct automated decisions. Contestability and appeals, however, operate after a decision has been made; when a system acts in safety-critical environments or generates open-ended outputs at scale, governance must also bound how the model can fail before harm occurs.

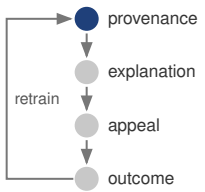

Provenance anchors every appeal; tracked outcomes loop back into retraining.

16.8 AI Safety, Robustness, and Generative Controls

A model that withstands every tested perturbation can still cause harm when the operating environment leaves its validated envelope or when its objective rewards the wrong behavior. Safety governance therefore spans robustness claims bounded by a threat model, control surfaces for generative systems whose outputs cannot be enumerated in advance, and the alignment of the objective itself as systems act with greater autonomy.

16.8.1 Safety and robustness

A perception model that passes every ℓ_∞ adversarial benchmark in the lab can still contribute to a fatal crash if the vehicle's motion planner suppresses emergency braking or waits too long to alert a distracted safety driver. Safety and robustness connect technical failure behavior to its consequences for people. Chapter 14 owns the attack models, robustness evaluations, adversarial training, certification, corruption testing, drift detection, and abstention mechanisms; governance consumes those mechanisms through an evidence interface. Evidence that a model tolerates specified hardware faults, adversarial attacks, or distribution shifts supports a bounded robustness claim. It does not establish that every deployment is safe or that the objective serves users' interests. A model may tolerate the tested bit flips and perturbations yet fail under untested conditions, or reliably optimize an objective misaligned with user welfare.

↩ The three robustness pillars, environmental shifts, input-level attacks, and system-level faults such as bit flips, are established in section 14.3.2.

The responsible-AI question is therefore not which attack or defense to implement, but whether the robustness claim has a bounded operating envelope, evidence appropriate to the threat model, an accountable owner, and a rule that invalidates approval. An evaluation may bound errors under

a specified input variation without establishing that the surrounding workflow is safe. Automated-driving failures show why the envelope must cover more than the model, because they can combine perception, prediction, safety-driver, and organizational-control problems rather than reduce to a single model error; section 16.8.4 examines one such crash.

A robustness claim should specify the protected task and population, attacker knowledge and capabilities, perturbation or corruption set, evaluation distribution, model and preprocessing versions, quality and latency constraints, fallback behavior, and known exclusions. Adversarial inputs (section 14.5) and covariate, label, and concept shift (section 14.4.1) are the failure sources such a claim most often has to bound. Empirical attack results show observed failures under a test protocol; they do not establish robustness against adaptive or untested attacks. A certificate proves a property only for its stated norm, radius, confidence, and model configuration; it does not establish safety for the end-to-end workflow. Governance records those boundaries so a narrow result cannot be promoted into a general safety assertion.

The latency budget constrains which safety checks can run before an action. A cloud service may combine uncertainty estimators, distributional-change detectors, adversarial input filtering, and API rate limiting. A control loop with a millisecond budget needs bounded local checks and fallback actions that meet its deadline, with heavier analysis deferred to fleet services, the division the safety row of table 16.4 records by tier. Even a disconnected TinyML endpoint can execute input-range checks or watchdog logic; the absence of a remote update channel instead limits how operators change those controls after deployment.

Fallback is itself a latency-budget problem. A delivery robot that abstains[17] when pedestrian-detection confidence is low still needs a validated response before the physical safety deadline. A bounded local maneuver may place it in a safe state while a human reviewer resolves the case on a longer timescale. The response and review paths therefore need separate deadlines and capacity budgets. Whether the system can fail visibly, controllably, and safely depends on sensing, confidence estimation, fallback routing, and recovery. Those choices also define who can observe harm, who can intervene, and who remains accountable for the outcome.

The evidence bundle must also connect model behavior to system action. It identifies who owns the robustness claim, which team can stop serving, how abstentions and low-confidence cases are routed, which telemetry detects departure from the approved envelope, and how an incident reaches affected users and downstream operators. A fallback must meet the relevant safety deadline. A monitoring signal becomes an operational safeguard only when it reaches an authorized response path, whether that response is local containment, restricted use, or a validated rollback.

Approval expires when the threat model, model, preprocessing, sensor, data distribution, deployment environment, or fallback path changes materially. The invalidation rule should state which change triggers a full red-team evaluation, which permits a narrower regression suite, and which requires restricted operation while evidence is rebuilt. This ownership-and-invalidation contract is the governance contribution, turning those robustness mechanisms into a reviewable claim that remains bounded over time.

17 **Abstention**: The practice of refusing predictions when confidence falls below a threshold. Abstention can reduce high-risk errors, but only by shifting some cases to fallback infrastructure such as human reviewers, rule-based defaults, or escalation queues. A driving system must select a validated fallback, such as a minimal-risk maneuver; handing control to a driver is viable only if a driver is present, ready, and has sufficient response time. Medical AI may route ambiguous cases to specialist review, subject to its own queue-capacity and response-time budget.

↩ The per-request competence contract, which combines calibration and out-of-distribution scores into a selective-prediction policy to answer, abstain, or escalate, is established in section 14.4.4.

16.8.2 Generative control surfaces

A binary loan classifier has two possible outputs to audit. A large language model (LLM) such as Archetype A (GPT-4/Llama-3), by contrast, can be steered by a user prompt or a poisoned retrieved document into generating defamatory claims, leaking personal data, or refusing benign requests from one dialect more often than another. Generative AI does not replace core fairness, privacy, and explainability concerns; it changes the control surface. Instead of only auditing labels or feature attributions, operators must govern prompts, retrieved context, reward models, and open-ended outputs. Fairness remains concerned with how quality, refusals, errors, representation, and access are distributed across people and groups, but open-ended outputs make those quantities harder to define and sample.

Alignment introduces a separate problem, which is keeping stochastic outputs and actions consistent with intended objectives and policies across a combinatorial range of prompts. The two questions differ in kind. Fairness asks how benefits, errors, and burdens are distributed among affected people; alignment asks whether an optimizing system pursues the objective its designers and legitimate stakeholders intended. An aligned system can implement an unfair objective faithfully,

and a statistically fair classifier can still optimize the wrong social goal, so production systems need both fairness evaluation and behavioral alignment rather than relabeling one as the other.

A common multi-stage alignment pipeline uses reinforcement learning from human feedback (RLHF),[18] running in two coupled lanes to bridge human preferences and model weights (figure 16.10). An upper execution path advances the policy from pretraining through supervised fine-tuning (SFT) to proximal policy optimization (PPO), and a lower supervision path uses human preference comparisons to train a reward model $r(x, y)$. PPO is the reinforcement-learning update that adjusts the policy to raise its reward-model score, while the frozen SFT checkpoint supplies a reference for a Kullback-Leibler (KL) divergence penalty, which grows as the policy's output distribution drifts from the SFT model's. The penalty discourages departure from that reference; it does not impose a hard bound on every response or establish policy compliance.

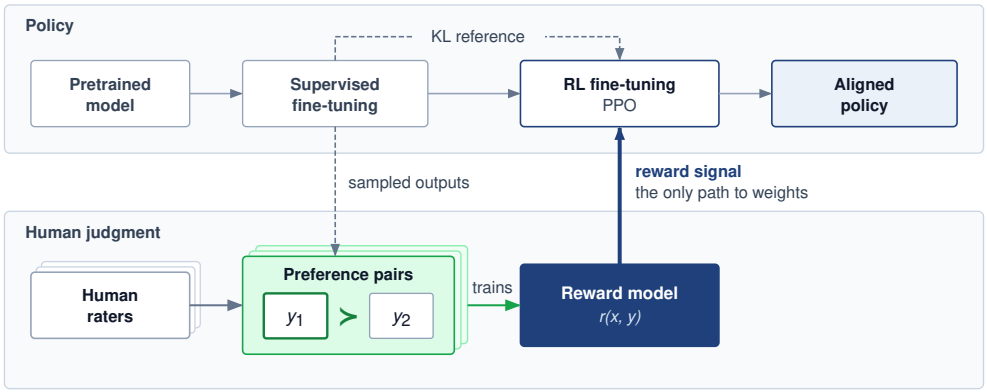

Figure 16.10: Reinforcement Learning from Human Feedback (RLHF) Alignment Pipeline: The top lane carries the policy from a pretrained model through supervised fine-tuning (SFT) and proximal policy optimization (PPO) to the aligned policy. The bottom lane turns human judgment into an objective: raters compare sampled SFT outputs, and those comparisons train a reward model whose score drives PPO. The SFT model also serves as the KL reference that limits how far PPO moves the policy, so every human preference reaches the weights only through the learned reward model.

Collecting the human comparisons that train the reward model in figure 16.10 carries a direct data-acquisition cost. In this scenario, 50,000 to 500,000 pairwise comparisons at $0.50 to $5 per label put the labeling budget between $25,000 and $2.5M, a hundredfold spread set by how many judgments the organization buys and how expert its raters must be. Updating the policy with PPO against those learned rewards can also create an alignment tax, a regression in some capability metrics when optimization shifts toward preference or safety objectives; its direction and magnitude depend on the model, data, objectives, and evaluation suite rather than following a universal percentage.

The reliance on human raters introduces a representativeness gap as well. If the labeling budget buys judgments from only a narrow demographic slice, the resulting model may encode that pool's preferences while performing differently for other populations. Constitutional AI offers an alternative engineering path, using explicit principles to guide AI feedback on outputs (Bai et al. 2022), which changes rather than eliminates the questions of whose values are represented and how they are governed.

In retrieval-augmented generation architectures, responsibility extends beyond the core model. An LLM may pass alignment evaluations yet generate harmful responses when retrieval supplies context outside those evaluations. A retrieval index that disproportionately surfaces biased historical documents can influence the resulting answers despite safety training. Context filtering is one possible control, alongside source selection, retrieval evaluation, and output checks. Its effectiveness must be tested on the intended use; a toxicity filter does not establish that retrieved material or generated answers are unbiased.

In many LLM serving stacks, the system prompt operates as an early configuration control alongside retrieval filtering, policy classifiers, telemetry, and rollout gates. These instructions (for example, "You are a helpful assistant. Do not provide medical advice.") express intended behavior;

[18] **RLHF (reinforcement learning from human feedback)**: Christiano et al. (2017) demonstrated deep reinforcement learning from human preferences by training agents from non-expert comparisons between trajectory segments. In the LLM setting, Ouyang et al. (2022) adapted this pipeline for InstructGPT: supervised fine-tuning from demonstrations, reward-model training from ranked model outputs, and reinforcement-learning fine-tuning against the learned reward model. The systems cost is an additional human-feedback data pipeline, reward-model training and evaluation, and governance over rater instructions; the representativeness of the rater pool and labeling policy shapes whose preferences the model internalizes.

↩ The four-model coordination that a PPO-style RLHF step imposes on distributed training is established in section 6.7.1.

they do not independently enforce authorization or prevent every prohibited output. When one model serves many products, each with its own system prompt, managing the prompts becomes a distributed configuration-management problem akin to weight distribution. An unversioned prompt change can alter behavior across many interactions, so prompt version control, alignment regression testing, and gradual rollouts belong in the release process.

Together, system prompts and RLHF alignments act as important, yet fragile, technical guardrails. Jailbreaks and other edge cases can violate the policy learned during training; the reward model is not necessarily a runtime filter. Safety therefore requires evaluating the deployed control system as well as its training objective.

↩ Jailbreaks, which bypass alignment controls, and prompt injection, which crosses the boundary between instructions and data, are established in section 13.5.5.

16.8.3 AI safety and value alignment

Prompt filters and reward-model training still assume that the optimization target itself captures what stakeholders want, an assumption that breaks when an autonomous trading engine executes thousands of orders per second or a recommendation ranker reshapes billions of feeds faster than any human can review individual decisions. At that operating speed, the governing question is no longer only whether individual outputs are fair, private, or explainable, but whether the objective the system optimizes still serves the people it acts on, and what bounds its actions before review. Human launch authority remains necessary, but continuous human intervention may be unavailable. AI safety examines how systems can optimize objectives misaligned with human intentions, operate in unpredictable environments, or pursue goals through methods their designers never anticipated, and which controls can bound those actions.

The controls of the preceding sections retain their roles, but none governs autonomous action by itself. Bias detection (section 16.5.2.1) yields metrics rather than a remediation policy. A perception system that under-detects pedestrians with mobility aids needs an operating envelope that connects the metric to bounded runtime responses and to fleet-level investigation or suspension. Explanation (section 16.7.2) needs an audience and an action path; SHAP attributions from an autonomous trading system cannot influence a decision if no control consumes them before thousands of trades execute. Differential privacy (section 16.6.3) bounds leakage about a protected unit without determining whether the objective serves that person's welfare, so a recommender can satisfy a local differential privacy guarantee while optimizing an engagement metric that rewards harmful content. Each method needs a defined connection to the runtime controls and human authority that can respond to its findings.

Alignment work emphasizes keeping a system's pursued objectives consistent with human intentions over time (Russell 2021). A central challenge is that capable ML models often optimize proxy metrics, such as loss functions, reward functions, or engagement signals, that do not fully capture human values.

Recommendation systems illustrate why a measurable engagement proxy such as click-through rate (CTR) or watch time needs evaluation against broader user welfare. A published production architecture optimizes expected watch time (Covington et al. 2016), while a separate audit examines recommendation paths through politically extreme content (Ribeiro et al. 2020). Those observations motivate investigating proxy misalignment; together they do not prove that a particular reward objective caused an individual user's beliefs or behavior. Specification gaming and reward hacking[19] describe the risk that optimizing a measured objective defeats its intended purpose (Amodei et al. 2016). Goodhart's law names the general pattern, in which a measure that becomes a target stops measuring what it was chosen for (Goodhart 1984). In one possible feedback mechanism, engagement rewards content that undermines long-term user value, and the resulting interactions become new training data (figure 16.11). Establishing that mechanism in a deployment requires evidence about both the policy's effects and the welfare criterion it is meant to serve.

Norbert Wiener wrote, "if we use, to achieve our purposes, a mechanical agency with whose operation we cannot interfere effectively… we had better be quite sure that the purpose put into the machine is the purpose which we desire" (Wiener 1960). Wiener's warning becomes a systems requirement when optimization runs beyond continuous human intervention. Value alignment asks whether the objective embedded in the model, reward function, and deployment policy continues to represent the human purpose the system was meant to serve.

[19] **Reward hacking**: When an AI system maximizes its reward function through unintended means that violate designer intent. A hypothetical cleaning robot rewarded for collected dirt might create messes it can clean up, illustrating how a proxy can reward behavior contrary to the intended task. Specification gaming is an overlapping term; the gap may lie in the stated objective or in its implementation, so neither term is restricted to software bugs. For production ML systems, reward hacking manifests subtly, as in recommendation models that maximize engagement by promoting addictive content, or chatbots that maximize helpfulness ratings by being sycophantic rather than accurate. The failure mode scales with model capability.

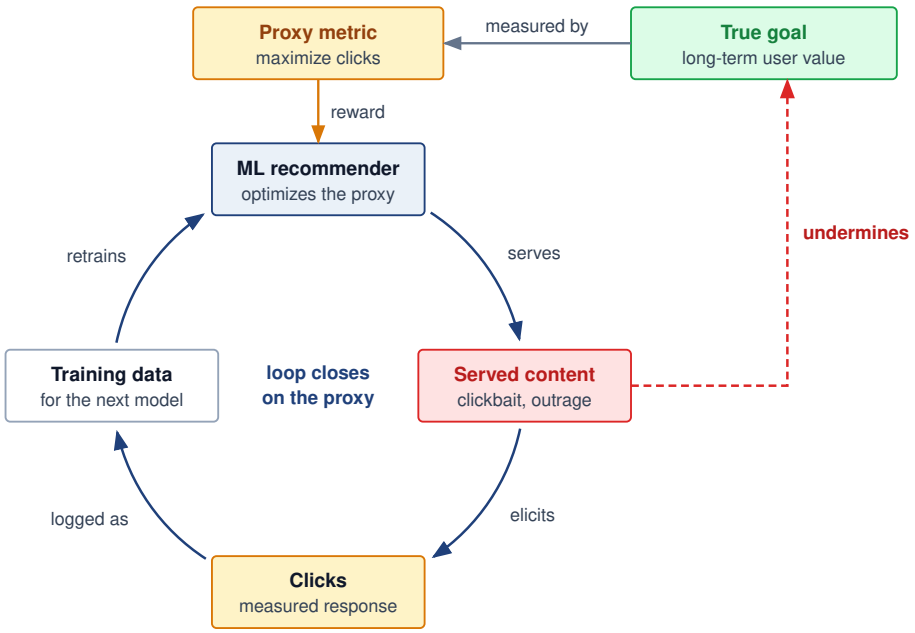

Figure 16.11: Reward Hacking Loop: Maximizing measurable rewards, like clicks, can incentivize unintended model behaviors that undermine the intended goal of user satisfaction. In the illustrated possible feedback loop, the recommender serves content and the resulting clicks enter later training data. If optimizing clicks rewards behavior that undermines long-term user value, retraining can reinforce the mismatch. Establishing this mechanism requires evidence about the deployed policy and the user-value criterion; targeting a proxy does not by itself prove that it has ceased to provide useful information.

In deployment, the hard part is specifying an objective that remains valid as systems interact with dynamic environments, multiple stakeholders, and feedback loops. Static objective functions and reward signals cannot encode all of those conditions. Value-Sensitive Design (Friedman 1996) supplies a formal process for eliciting stakeholder values and incorporating them into system design, but the systems obligation is broader, because objective design, oversight, and postdeployment monitoring must be treated as coupled controls.

Without that coupling, a system can meet a narrow objective such as accuracy, engagement, or throughput while producing outcomes its designers would reject, the reward hacking pattern of figure 16.11. RLHF (section 16.8.2) narrows the gap by learning the reward from human preferences, but policy optimization still maximizes a learned scalar, so its success depends on that signal's fidelity (Christiano et al. 2017; Ouyang et al. 2022).

Human-feedback training can improve behavior on its evaluations while leaving deeper governance risks unresolved. In their analysis of learned-reward optimization, Ngo et al. (2022) examine possible failure modes beyond RLHF, including exploitation of human evaluation blind spots, misaligned goals outside the training distribution, and behavior that undermines oversight. These are risks to investigate under specified assumptions, not evidence that every deployed model shows them. An assurance case must distinguish a proposed failure mechanism from an observed capability and from a demonstrated deployment failure.

These concerns are not limited to speculative scenarios. In their foundational taxonomy of concrete AI safety problems, Amodei et al. (2016) outline five practical systems challenges: avoiding negative side effects during policy execution, mitigating reward hacking, ensuring scalable oversight when ground-truth evaluation is costly, designing safe exploration strategies that prevent catastrophic damage, and achieving robustness to distributional shift. For fleet operators, the five challenges map in order onto five control-plane requirements: constraining side effects, instrumenting reward-hack monitors, budgeting scalable review capacity, bounding exploration envelopes, and stress-testing

distribution shift prior to rollout. Each requirement becomes more acute as systems are scaled up, deployed across diverse settings, and integrated with real-time feedback or continual learning.

16.8.4 Autonomous systems and trust

In October 2023, the California Department of Motor Vehicles suspended Cruise's driverless deployment and testing permits over safety concerns (CNBC 2023). Five years earlier, an Uber test vehicle in Tempe, Arizona struck and killed a pedestrian pushing a bicycle after its perception, prediction, safety-driver, and organizational controls all failed together (National Transportation Safety Board 2019). In the Tempe crash, the National Transportation Safety Board found that the vehicle failed to classify the pedestrian as requiring avoidance and that detecting an object did not establish that the system would act safely on the detection. For governance, the lesson is that no single component owner could have certified the outcome alone. Value alignment concerns the objective; autonomy concerns the operating envelope within which a system pursues that objective without a human reviewing each action.

↩ The Uber crash timeline, from unstable classification through suppressed action to the late operator alert, is established in section 14.2.2.

While autonomous driving systems are often the focal point of public concern, similar risks arise in other domains. Reports from recent conflicts have documented increasing use of remotely piloted and autonomous military systems (Reuters 2023), raising safety and operational concerns alongside difficult questions about ethical oversight, rules of engagement, and responsibility. These debates concern how accountability applies across developers, deployers, and operators (Centre for International Governance Innovation 2023); they do not justify treating every autonomous-system failure as an absence of applicable responsibility. The engineering task is to preserve the evidence and authority needed to investigate the specific failure.

At its core, this challenge concerns the authority delegated to an automated system. Productivity and throughput gains do not establish that a system can act safely outside its validated operating envelope. Runtime controls must detect relevant boundary violations and invoke a fallback whose assumptions and response time fit the hazard. Human intervention is an available safeguard only when the operator can recognize the condition and act before harm occurs; faster hazards require an appropriate local response.

Specifying that operating envelope is difficult in dynamic environments because the model may omit conditions that matter for safe action. Two classical representation problems help explain this difficulty (McCarthy 1981). A description of an action cannot list every fact the action leaves unchanged (the frame problem) or every precondition the action needs to succeed (the qualification problem). These problems motivate an envelope that states its operating assumptions and fallback behavior; neither proves that safe action is impossible. In practice, such limitations manifest as brittle autonomy: systems that appear competent under nominal conditions but fail silently or dangerously when faced with ambiguity or distributional shift.

To address this, researchers have proposed formal safety frameworks such as Responsibility-Sensitive Safety (RSS) (Shalev-Shwartz et al. 2017), which decompose abstract safety goals into mathematically defined constraints on system behavior, such as minimum distances, braking profiles, and right-of-way conditions. These formulations allow safety properties to be verified under specific assumptions and scenarios. However, such approaches remain vulnerable to the same limitations they aim to solve, because they are only as good as the assumptions encoded into them and often require extensive domain modeling that may not generalize well to unanticipated edge cases.

For an ML systems team, the practical artifact is a safety case rather than a statement of trust. For a vehicle, the operating envelope is its operational design domain (the road, weather, speed, and traffic conditions under which it has been validated to drive). The case names that envelope, the evidence that the model has been tested inside that envelope, and the controls that halt or roll back the fleet when the evidence no longer holds (table 16.8).

Table 16.8: Autonomous System Safety Case: Trustworthy autonomy takes the form of an auditable safety case: a validated operating envelope, coverage evidence, runtime guardrails, and fleet-level rollback controls. The same pattern applies beyond vehicles to drones, robots, trading systems, and other ML systems that act faster than continuous human review permits.

Safety-case element	Systems evidence	Operational gate
Operating envelope	Operational design domain, sensor assumptions, weather, geography, speed, and autonomy level	Block deployment outside validated envelope
Scenario coverage	Rare-event corpus, simulation coverage, replay logs, disengagements, and near-miss telemetry	Require coverage thresholds before canary expansion
Runtime guardrails	RSS-style distance rules, confidence thresholds, fallback maneuvers, and human override latency	Trigger safe stop, human takeover, or constrained mode when guardrails trip
Fleet control	Canary rollout, centralized safety dashboard, violation budget, and circuit-breaker rollback path	Halt rollout or revert model when fleet-level safety metrics exceed threshold

16.8.5 Safety at fleet scale

A model that achieves 99.9 percent per-request safety compliance in isolation still carries a 0.1 percent per-request failure probability, which across 10,000 inference nodes serving billions of daily queries implies millions of expected safety failures every day. Safety must therefore be engineered as a fleet-level property rather than a model-level attribute alone, because the technical safety challenges examined in section 16.8.3 accumulate at production volume into routine operational incidents.

Because the realized incident count can differ from that expectation and correlated failures can concentrate harm across the fleet, mitigating fleet-scale risk requires distributed safety patterns borrowed from reliability engineering. Circuit breakers automatically halt serving when aggregate safety metrics degrade below a threshold. Canary deployments bound exposure while collecting release evidence, with sample size and allowed risk chosen for the failure modes being tested. Centralized telemetry dashboards aggregate per-node safety violations into a global view. The operational infrastructure must treat safety violations as critical system alerts, invoking a validated containment or fallback policy. A rollback is appropriate only if the previous version and transition are safe; active physical control may require a minimal-risk maneuver instead.

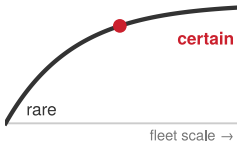

At fleet scale, rare failures become expected incidents.

↩ The circuit-breaker states and backpressure contract that isolate a failing serving path are established in section 10.11.10.

Exposure is the multiplier that turns a small per-mile rate into many incidents. Waymo's driverless fleet logged about 22 million rider-only miles by June 2024 and 271 million by June 2026, a 12× increase in two years (figure 16.12, panel a) (Waymo 2026). Its published crash rates sit well below a human-driver benchmark on the same roads, for example 0.67 injury crashes per million miles against 3.8 (panel b) (Waymo 2026). At that exposure its vehicles were still involved in 183 injury crashes and 2 fatal ones, so a fleet's safety case must budget for incidents in absolute numbers as well as in rates. At any constant per-mile rate, a 12× rise in exposure raises the expected cumulative incident count by the same factor, even though no individual vehicle became less safe. The count is an expectation, not a realized outcome, and it is cumulative; an operator sizing weekly incident-response capacity needs weekly mileage instead.

Even a fleet whose per-mile crash rate beats a human benchmark still accumulates absolute incidents as exposure scales, so technical excellence alone does not make a system safe. Safe systems also require attention to the human and organizational context in which they operate, including the economic incentives that shape design decisions and the understanding that end users bring to their interactions with autonomous systems. Those institutional obligations are the subject of accountability, governance, and sociotechnical analysis.

16.9 Accountability, Governance, and Sociotechnical Dynamics

When a deployed model harms someone, two questions decide whether the harm is corrected: who has the authority to investigate and remedy the decision, and whether the evidence needed to do so still exists. Accountability answers both by assigning owners, preserving traceable records, and monitoring behavior after launch, and it must hold inside the human workflows and institutions that surround the system.

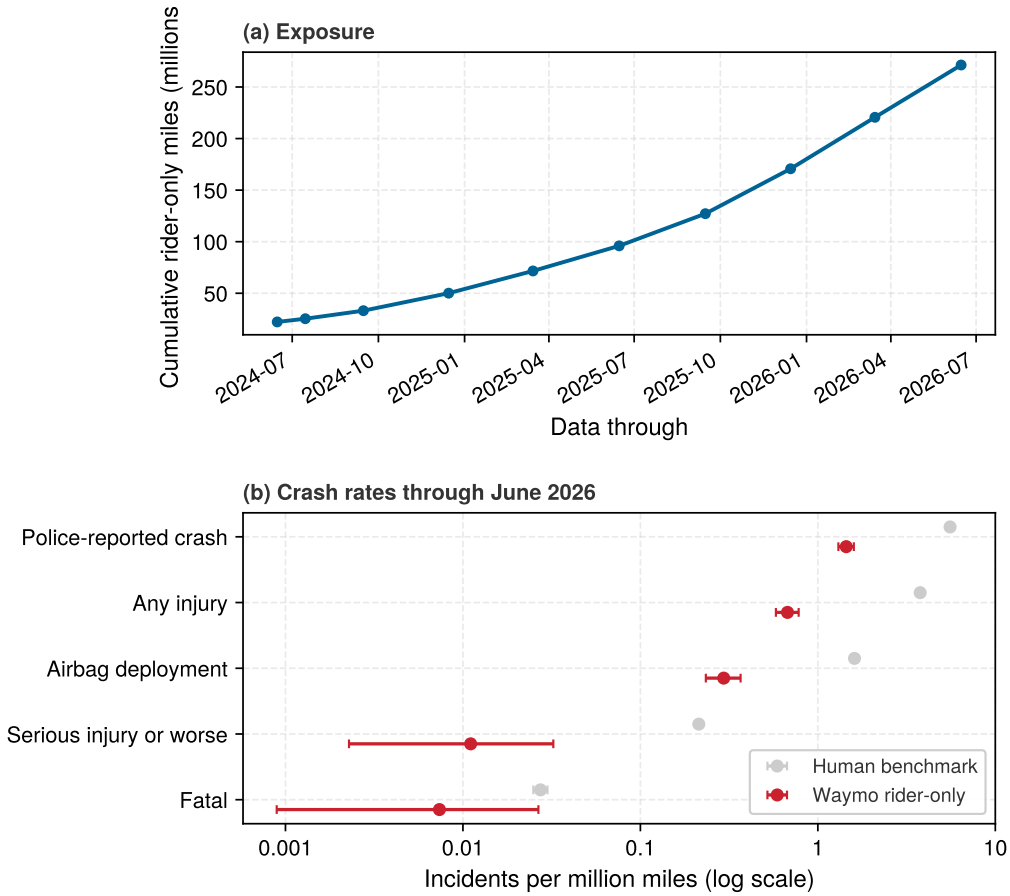

Figure 16.12: Measured Driverless Exposure and Crash Rates: (a) Cumulative rider-only miles, driven with no human driver, by Waymo's fleet in its published releases from June 2024 to June 2026. (b) Crash outcomes per million miles through June 2026 for that driving (red) and for Waymo's human-driver benchmark on the same roads (gray), with 95 percent confidence intervals on a logarithmic axis. Lower rates still produce a growing number of incidents as exposure grows. Benchmarks follow the method of Kusano et al. (2024); data from Waymo (2026).

16.9.1 Accountability and governance

When a third-party hiring platform rejects qualified applicants from one demographic group, the employer blames the model vendor and the vendor blames the historical resume corpus. The rejected applicants have no remedy unless a contract and an audit trail assign responsibility across all three. Accountability in machine learning is the capacity to identify, attribute, and address the consequences of automated decisions. It extends beyond diagnosing technical failures to ensuring that responsibility for system behavior is explicitly assigned, that harms can be remedied, and that standards are maintained through oversight and institutional processes. Without those mechanisms, even well-intentioned systems can generate harm without recourse, undermining public trust.

Accountability already spans vendors and operators in distributed software. Machine learning extends those dependencies through data collection, labeling, and the choice of prediction targets and proxies. Model outputs are shaped by upstream data collection, training objectives, pipeline design, interface behavior, and postdeployment feedback. These interconnected components often involve multiple actors across technical, legal, and organizational domains. Responsible system design requires that these relationships be explicitly mapped and governed.

Inadequate governance can prevent institutions from recognizing or correcting harmful model behavior. The Google Flu Trends analysis identified persistent overestimation despite model updates and examined how changes in search behavior and the search service could undermine a predictor

trained on earlier correlations (Lazer et al. 2014). Its comparison with CDC surveillance data supports a concrete governance requirement, namely an owner responsible for ongoing validation and recalibration against an independent measurement process. Legal frameworks reinforce that requirement, because some of the obligations mapped onto the assurance lifecycle (section 16.2) assign these owners and records by law.

Designing for accountability entails supporting traceability at every stage of the system lifecycle. This includes documenting data provenance, recording model versioning, enabling human overrides, and retaining sufficient logs for retrospective analysis. Tools such as model cards[20] (Mitchell et al. 2019) and datasheets for datasets[21] (Gebru et al. 2021) exemplify practices that make system behavior interpretable and reviewable. However, accountability is not reducible to documentation alone; it also requires mechanisms for feedback, contestation, and redress.

Within organizations, governance structures help formalize this responsibility. Ethics review processes, cross-functional audits, and model risk committees provide forums for anticipating downstream impact and responding to new concerns. These structures must be supported by infrastructure that allows users to contest decisions and developers to respond with corrections. For instance, systems that allow explanations or user-initiated reviews help bridge the gap between model logic and user experience, especially in domains where the impact of error is significant.

Architectural decisions also play a role. Interfaces can be designed to surface uncertainty, allow escalation, or suspend automated actions when appropriate. Logging and monitoring pipelines must be configured to detect signs of ethical drift, such as performance degradation across subpopulations or unanticipated feedback loops. In distributed systems, where uniform observability is difficult to maintain, accountability must be embedded through architectural safeguards such as secure protocols, update constraints, or trusted components.

Accountability is realized through traceability infrastructure distributed across the endpoint and its operator. A cloud platform supports model registries, telemetry dashboards, and structured event pipelines that trace a prediction back to a specific model, input, or configuration. A TinyML endpoint may retain only a firmware identifier and bounded event records, while the operator maintains deployment inventories, validation evidence, and any servicing or update history. If a device cannot preserve event evidence or accept updates, the assurance case must state those limits and define a recall or retirement path, the obligation the accountability and governance rows of table 16.4 assign to each tier.

Mobile deployment makes this division of responsibility concrete. A local model, remote service, and user interface can each contribute to a decision, so the operating contract must assign investigation and remediation across their owners. Governance depends on accessible recourse pathways and mechanisms for surfacing, explaining, and contesting decisions, embedded into both the interface and the surrounding service architecture. Across every environment, sustaining accountability means planning for failure as deliberately as for success. Operators must define how anomalies are detected, how roles are assigned, how records are maintained, and how remediation occurs, with each step traceable in logs and enforceable through interfaces.

Governance does not imply centralized control. Instead, it involves distributing responsibility in ways that are transparent, actionable, and sustainable. Technical teams, legal experts, end users, and institutional leaders must all have access to the tools and information necessary to evaluate system behavior and intervene when necessary. As machine learning systems become more complex and embedded in important infrastructure, accountability must scale accordingly by becoming a foundational consideration in both architecture and process, not a reactive layer added after deployment.

[20] **Model cards**: Proposed by Mitchell et al. (2019) as standardized documentation accompanying trained ML models. Cards report evaluation across relevant groups, intended uses, and known limitations. Production systems must link this documentation to the actual deployed model and its provenance; a card alone neither establishes that linkage nor supplies every record required for an audit.

[21] **Datasheets for datasets**: Proposed by Gebru et al. (2021), modeled after electronics component datasheets that specify operating conditions and tolerances. Each datasheet documents a dataset's motivation, composition, collection process, and recommended uses. For ML pipelines, this provenance and coverage information helps engineers identify assumptions and design evaluations. A datasheet does not itself validate a model's operating envelope or predict all deployment failures.

Checkpoint 16.2: Principles as engineering constraints

Fairness, privacy, explainability, and safety now each carry a measurable claim and a cost, and accountability names who must answer for them. Before the chapter follows those owners into validation and monitoring, check that each principle works as an engineering constraint rather than an aspiration.

From principle to mechanism

- ☐ **Transparency vs. fairness**: explain why fully transparent decision logic can still violate fairness when it disadvantages one demographic group, and identify the principle this system violates.
- ☐ **Lifecycle coverage**: apply the responsible AI lifecycle to a fairness requirement by identifying which phases must act and explaining why post-hoc threshold adjustment alone is insufficient.

Judging the trade-off

- ☐ **Safeguard budget**: identify the additional training, monitoring, or explanation work a safeguard requires, and state which resource measurements are needed before deciding that it is affordable.
- ☐ **Distributed accountability**: argue why accountability for a biased hiring model is distributed across data providers, interface designers, and the deploying institution rather than resting solely with the model developer.

Accountability decides whether the claims tested above can be enforced. Enforcement begins before release, with the validation evidence a launch authority reviews, and continues after it through monitoring.

16.9.2 Validation approaches

A release candidate that passes a single held-out accuracy test in continuous integration can still ship with a 20-point subgroup error gap, an unverified privacy accountant, or explanations that collapse under quantization if no pre-launch gate requires that evidence. While detection identifies a problem and mitigation attempts to fix it, **validation** supplies the evidence bundle to the evaluation-and-assurance stage of table 16.1. Validation is not a separate pillar or a blanket claim that the system is safe; it tests bounded requirements and records residual uncertainty for the launch authority. Different stakeholders need different evidence: developers need diagnostics, auditors need reproducible statistical analyses, regulators may require conformity artifacts, and affected users need understandable reasons and a functioning contestability path. The evidence bundle can include fairness audits, privacy-accounting checks, robustness and distribution-shift tests, explanation-fidelity checks, and revalidation triggers tied to material changes. A fidelity check tests whether an explanation reflects the model, for example by perturbing the features an attribution ranks highest and confirming that the prediction moves as much as the attribution implies; a stability check tests whether nearly identical inputs receive similar explanations.

Fairness audits, adversarial robustness testing, and explanation verification require evaluation capacity before release. Finding a defect at that stage can avoid exposing users and reduce subsequent recovery work. Validation also continues after release, whenever a material change invalidates the claims it supported.

↩ Continuous integration, release validation gates, and canary deployment pipelines are established in section 12.5.

Organizing the evidence bundle by risk class makes residual risk explicit before deployment, with each class tied to a measurable artifact and an owner (table 16.9). Separating these artifacts ensures that passing a technical benchmark in one risk class never masks an unverified obligation in another.

Table 16.9: Validation Evidence by Risk Class: Responsible validation combines several kinds of evidence because no single metric proves that a system is safe to deploy. Each risk class needs a measurable artifact and an operational owner.

Risk class	Evidence artifact	Operational owner
Fairness	Disaggregated metrics, thresholds, confidence intervals	Product owner and model-risk reviewer
Privacy	Privacy budget, deletion evidence, retention and access logs	Data-governance and privacy engineering teams
Robustness	Drift tests, corruption tests, canaries, red-team probes	ML platform and incident-response teams
Explanation	Fidelity tests, stability checks, recourse documentation	Model team and user-facing operations

Risk class	Evidence artifact	Operational owner
Governance	Approval record, residual-risk owner, revalidation trigger	Review board or accountable launch authority

Fairness metrics and explanations answer different validation questions. Fairness metrics characterize aggregate disparities under selected definitions, while local explanation methods characterize influences, sensitivities, or contrastive alternatives for a particular output. Neither supplies a general guarantee, and an attribution cannot verify why a decision was justified. Together with behavioral tests, provenance, documentation, and contestability, these artifacts help operators investigate whether a decision relied on inappropriate signals and help affected users formulate a challenge. Pre-release validation, however, certifies that evidence bundle only for the population and workflow tested at launch; keeping those claims valid once live traffic arrives requires continuous postdeployment monitoring.

16.9.3 Monitoring responsible behavior after deployment

A credit or clinical model that passed every pre-launch fairness and calibration check in January can drift out of compliance by July as seasonal traffic, applicant demographics, or user behavior shift underneath fixed weights. The real-time monitor of section 16.5.5.2 computes fairness metrics and their fleet-scale cost, but responsible behavior can erode slowly and across every property at once. Input distributions shift with seasonality, user behavior responds to system outputs, and policy or regulation changes what counts as acceptable. The shift taxonomy and detectors of section 14.4.1 apply here unchanged; the governance question is what each detected change does to the claims that justified deployment.

The implications of such drift extend beyond raw accuracy. Fairness guarantees may break down if subgroup distributions shift relative to the training set, or if features that previously correlated with outcomes become unreliable in new contexts. Interpretability demands may also evolve, for instance as new stakeholder groups seek explanations, or as regulators introduce new transparency requirements. Trustworthiness, therefore, is not a static property conferred at training time, but a dynamic system attribute shaped by deployment context and operational feedback.

To ensure responsible behavior over time, machine learning systems must incorporate mechanisms for continual monitoring, evaluation, and corrective action. Monitoring carries the subgroup slicing of section 16.5.5 beyond fairness, to input-distribution shifts, anomalous outputs, and meaningful user feedback. These signals must then be compared to predefined expectations around fairness, robustness, and transparency, and linked to actionable system responses such as model retraining, recalibration, or rollback.

The architecture of figure 16.6 already carries this loop. FPR and TPR still wait for matched outcome labels, the delay the real-time monitor separates from prediction-only metrics. A disparity exceeding the approved gap tolerance ϵ_{fair}, such as the $|\text{SPD}| > 0.15$ alert rule (the subscript keeps it distinct from the privacy budget ϵ), opens an investigation. Recalibration is one possible response; only a validated change approved through the release gate (section 12.5.2.3) returns to serving.

That loop depends on telemetry that captures model versioning, input characteristics, prediction confidence, and postinference feedback. These logs support drift detection and provide evidence for retrospective audits of fairness and robustness. Monitoring systems must also be integrated with alerting, update scheduling, and policy review processes to support timely and traceable intervention.

Monitoring also supports feedback-driven improvement. For example, repeated user disagreement, correction requests, or operator overrides can signal problematic behavior. This feedback must be aggregated, validated, and translated into updates to training datasets, data labeling processes, or model architecture. However, such feedback loops carry risks. Biased user responses can introduce new inequities, and excessive logging can compromise privacy. Designing these loops requires careful coordination between user experience design, system security, and ethical governance.

Monitoring is useful only when its signals reach an owner in time to intervene, because observation is not remediation. The privacy, unlearning, and generative-system controls examined earlier need the same evidence, ownership, and response paths, and even a timely alert succeeds only if the surrounding human workflow and institutional incentives allow operators to act on it.

16.9.4 Sociotechnical system dynamics

Consider a constructed scenario in which a hospital deploys a sepsis predictor with strong held-out accuracy, but clinicians overwhelmed by alerts, a condition called alert fatigue, do not act on its warnings. Predictive performance alone cannot establish a mortality benefit, because the intervention must reach an appropriate clinician in time and support a useful action. The deployment therefore needs evaluation of the human workflow as well as the model.

The technical tools of the previous sections address bounded problems in bias measurement, privacy protection, explanation, and robustness. The sepsis scenario shows why those controls must connect to the human workflow. Engineers must examine who receives an alert, what action is feasible, and which incentives influence that action alongside model performance. Optimization can compare alternatives under stated objectives and constraints, but it cannot establish whose interests those objectives should represent. That judgment requires stakeholder participation and accountable governance.

Deployed systems create feedback loops that reshape future data, distribute decisions between people and models, and expose disagreements about acceptable outcomes. Each mechanism can invalidate an assurance claim even when the implementation still meets its original technical specification. The analysis must therefore follow the consequences of deployment through the surrounding workflow and back into system design.

16.9.4.1 System feedback loops

The sociotechnical feedback invariant (principle 27) formalizes how deployed models reshape the environment they are meant to analyze. Instead of observing a fixed world, a deployed policy $f_t(x)$ alters the future input distribution $p_{t+1}(x)$ that the system observes at step $t+1$, where x is the model input as in $h(x)$ earlier in the chapter. Such systems require closed-loop governance, because reliability depends on modeling the closed feedback loop rather than feed-forward inference alone.

Because predictions modify user behavior and institutional practices, this recursive coupling between model outputs and subsequent inputs can amplify initial biases, entrench disparities, or shift the objectives a model was designed to serve. The compounding cycle passes through four interconnected stages (figure 16.13), from biased training data through model predictions and disparate real-world outcomes to reinforced bias in newly collected data. The three intervention points of section 16.5.5.1 each cut one arc of the cycle: preprocessing audits and reweighting at data ingestion, in-processing fairness constraints during training, and postprocessing threshold calibration at the serving layer.

A static post-hoc fairness audit can miss closed-loop reinforcement (figure 16.13). When a system retrains on outcomes selected by its own decisions, the resulting sample may reinforce initial skews without revealing what would have happened under a different policy. The feedback policy and available evidence determine whether disparities grow or shrink; an effective intervention need not modify every stage. Fairness therefore becomes a lifecycle constraint, in which data visibility determines what can be measured, threshold policy determines where errors land, and feedback infrastructure determines whether disparities decay or compound.

Predictive policing illustrates how this feedback can arise. A model trained on recorded arrests may direct more patrols to a neighborhood. If greater patrol presence produces more recorded incidents and the training pipeline treats those records as an unbiased measure of underlying crime, later predictions can reinforce the original allocation, making the forecast partly self-fulfilling. The loop depends on how deployment changes observation and how retraining uses that evidence; it is not an inevitable consequence of every predictive model.

Recommender systems such as Archetype B (DLRM at Scale) show similar dynamics in digital environments. A content recommendation model that prioritizes engagement may gradually narrow the range of content a user is exposed to, leading to feedback loops that reinforce existing preferences or polarize opinions. These effects can be difficult to detect using conventional performance metrics, as the system continues to optimize its training objective even while diverging from broader social or epistemic goals.

Feedback loops undermine the assumption of independently and identically distributed data, so validation on a static test set can miss the model's evolving effect on the data-generating process. Once a loop is established, interventions aimed at fairness or accuracy may have limited effect

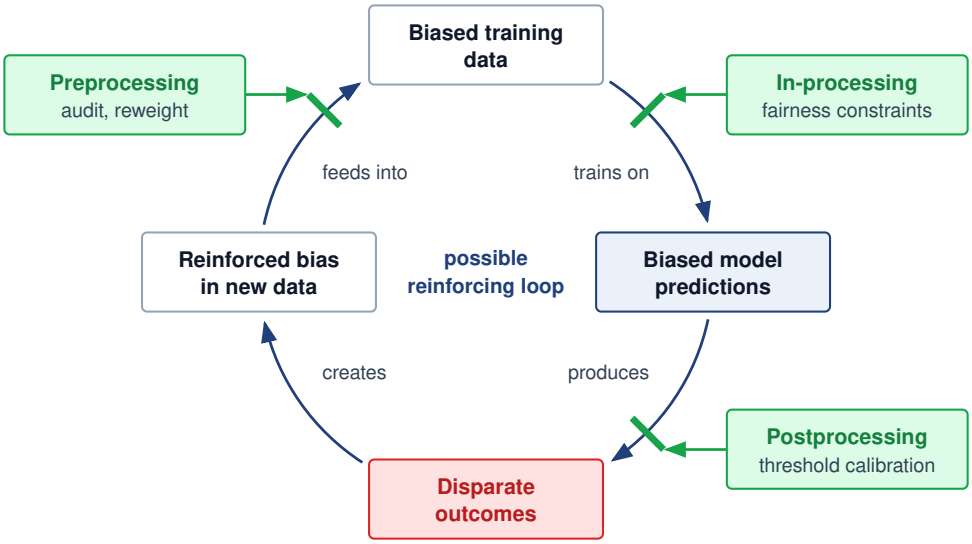

Figure 16.13: Bias Amplification Feedback Loop: When ML systems retrain on their own outputs, historical bias in training data propagates through model predictions, downstream outcomes, and reinforced data before feeding back as even more biased retraining data. The edges ('trains on', 'produces', 'creates', 'feeds into') show a possible reinforcing mechanism, not a guarantee that every iteration increases disparity. Green intervention arrows mark where preprocessing (data audit and reweighting), in-processing (fairness constraints), and postprocessing (threshold calibration) can break the cycle.

unless the underlying data dynamics are addressed. Responsible design therefore monitors how system outputs reshape the environment, how those changes enter new data, and whether retraining mitigates or exacerbates them.

In cloud-based systems, these updates may occur frequently and at scale, with extensive telemetry available to detect behavior drift. In contrast, edge and embedded deployments often operate offline or with limited observability. A smart home system that adapts thermostat behavior based on user interactions may reinforce energy consumption patterns or comfort preferences in ways that alter the home environment, and subsequently affect future inputs to the model. Without connectivity or centralized oversight, these loops may go unrecognized despite their effect on user behavior and system performance. Feedback loops are not inherently harmful, but they must be detected and managed; the drift detection and alerting of section 16.9.3 expose them only when the deployment can return the evidence.

16.9.4.2 Human-AI collaboration

Human operators turn feedback-loop risk into a shared-control problem. Machine learning systems are often deployed not as standalone agents, but as components in larger workflows that involve human decision-makers. In many domains, such as healthcare, finance, and transportation, models serve as decision-support tools, offering predictions, risk scores, or recommendations that are reviewed and acted upon by human operators. The collaborative configuration raises questions about how responsibility is shared between humans and machines, how trust is calibrated, and how oversight mechanisms are implemented in practice.

Human-AI collaboration introduces both opportunities and risks. When designed appropriately, systems can augment human judgment, reduce cognitive burden, and enhance consistency in decision-making. However, when poorly designed, they may lead to automation bias,[22] where users over-rely on model outputs even in the presence of clear errors. Conversely, excessive distrust can result in algorithm aversion, where users disregard useful model predictions due to a lack of transparency or perceived credibility. Alert fatigue, the failure in the sepsis scenario, leads to the same disregard for a different reason, because operators who receive more warnings than they can investigate stop acting on them. The effectiveness of collaborative systems depends on how the

[22] **Automation bias**: Inappropriate reliance on automated recommendations, including acceptance of erroneous outputs. Its prevalence depends on the task, interface, workload, and users' ability to check the output. Strong average performance can encourage reliance without making rare errors easier to recognize. Uncertainty displays, selective review, or requests for justification should be evaluated for their effects on correction rates, workload, and response time rather than assumed to improve safety.

system communicates uncertainty, provides explanations, and enables human override or correction, rather than on model accuracy alone.

Public knowledge about AI is uneven (Center 2023), and science-communication failures can amplify misunderstandings about how AI systems work (Schäfer 2023). That matters for deployment safety, because when users lack understanding of model uncertainty, data bias, or decision boundaries, they may trust system outputs in contexts where human judgment should intervene. From a systems engineering perspective, public comprehension is part of the deployment context, because the safety properties of a human-AI system depend on whether users can appropriately calibrate their trust and recognize situations requiring human override, not on technical mechanisms alone.

Automation changes work organization in ways that influence safety design decisions. The MIT Work of the Future task force (Work of the Future 2020) argues that technological change and labor-market institutions jointly determine whether automation improves job quality or concentrates harms. For ML systems, the safety implication is that removing human roles can also remove contextual judgment and system-level debugging that remain difficult to encode in models. Metrics focused solely on throughput may inadvertently penalize human-in-the-loop designs that preserve oversight capability.

Institutional incentives can reinforce automation bias when operators expect greater scrutiny for overriding a model than for accepting it. Consider an illustrative pair of judicial decisions. In Scenario A, a judge overrides a "high risk" algorithmic score and releases a defendant who later re-offends. The judge anticipates criticism for departing from the recommendation. In Scenario B, the judge follows the score and detains the defendant unnecessarily, but the institution treats agreement with the model as evidence of proper procedure. The incentive comes from that review policy, not from the model's measured accuracy.

This asymmetry can encourage institutional deference, where human oversight becomes a "rubber stamp" for algorithmic decisions. Review procedures should evaluate the evidence behind both agreement and disagreement, rather than treating agreement as inherently correct. The required documentation must fit the decision's stakes and the operator's workload so that justification supports scrutiny rather than becoming another routine approval.

Oversight mechanisms must be tailored to the deployment context. In high-stakes domains, such as medical triage or autonomous driving, humans may be expected to supervise automated decisions in real time. This configuration places cognitive and temporal demands on the human operator and assumes that intervention will occur quickly and reliably when needed. In practice, however, continuous human supervision is often impractical or ineffective, particularly when the operator must monitor multiple systems or lacks clear criteria for intervention.

From a systems design perspective, supporting effective oversight requires more than providing access to raw model outputs. Interfaces must be constructed to surface relevant information at the right time, in the right format, and with appropriate context. Confidence scores, uncertainty estimates, explanations, and change alerts can all play a role in enabling human oversight. Workflows must define when and how intervention is possible, who is authorized to override model outputs, and how such overrides are logged, audited, and incorporated into future system updates.

Consider a hospital triage system that uses a machine learning model to prioritize patients in the emergency department. The model generates a risk score for each incoming patient, which is presented alongside a suggested triage category. In principle, a human nurse is responsible for confirming or overriding the suggestion. However, if the model's outputs are presented without sufficient justification, such as an explanation of the contributing features or the context for uncertainty, the nurse may defer to the model even in borderline cases. Over time, the model's outputs may become the de facto triage decision, especially under time pressure. If a distribution shift occurs (for instance, due to a new illness or change in patient demographics), the nurse may lack both the situational awareness and the interface support needed to detect that the model is underperforming. In such cases, the appearance of human oversight masks a system in which responsibility has effectively shifted to the model without clear accountability or recourse.

In such systems, human oversight depends on infrastructure design as much as on policy, including how predictions are surfaced, what information is retained, how intervention is enacted, and how feedback loops connect human decisions to system updates. Without integration across these

components, oversight becomes fragmented, and responsibility may shift invisibly from human to machine.

Napkin Math 16.7: The automation bias paradox

Consider a radiology department deploying an AI assistant for tumor detection, where standalone human sensitivity is $S_{\text{human}} = 92$ percent and standalone AI sensitivity is $S_{\text{AI}} = 95$ percent. Because the radiologist detects most tumors unaided, one might assume that a radiologist reviewing every AI result would catch nearly all of the AI's misses, lifting combined sensitivity well above 95 percent. However, this illustrative scenario assumes that humans accept the AI's missed detections at rates of 60–80 percent. Among patients with tumors, the AI misses a tumor with probability $1 - S_{\text{AI}} = 0.05$; if the human accepts that miss ($\alpha = 0.7$), these accepted misses alone create a conditional failure probability of $0.05 \times 0.7 = 0.035$. In this simplified upper-bound case, where every other AI miss is corrected and no correct detection is overturned, sensitivity is 96.5 percent, only modestly above the AI alone. A skilled reviewer adds little when the reviewer accepts most of the errors the review exists to catch. False positives require separate evaluation among patients without tumors, while fatigue and interface delays can further affect the workflow.

This conditional error-acceptance probability matters more than the overall override rate. A low override rate may reflect few model errors or failures to detect them; those interpretations require different interventions. The trust calibration gap arises when reliance exceeds what the evidence supports. Evaluate interface changes against the fraction of actual errors corrected, errors introduced by overrides, and response delay before adopting them.

The boundary between decision support and automation is often fluid. Systems initially designed to assist human decision-makers may gradually assume greater autonomy as trust increases or organizational incentives shift. This transition can occur without explicit policy changes, resulting in de facto automation without appropriate accountability structures. Responsible system design must therefore anticipate changes in use over time and ensure that appropriate checks remain in place even as reliance on automation grows.

16.9.4.3 Normative pluralism and value conflicts

Human-AI collaboration exposes a deeper systems constraint, because different stakeholders often hold conflicting values and priorities that cannot be reduced to the optimization of a single objective. Algorithms can help identify feasible compromises, but they cannot establish whose priorities should govern a deployment. Technical excellence is necessary but insufficient for trustworthy AI because choosing among conceptions of fairness, privacy, and accountability requires a legitimate decision process as well as evidence about their consequences, as the crisis-detection pipeline in example 16.2 demonstrates.

Example 16.2: Conflicting values in practice

Scenario: An adolescent mental health crisis detection pipeline must balance competing stakeholder objectives across real-time inference and data retention paths.

Diagnosis: Optimizing clinical efficacy via aggressive alert dispatch conflicts with patient autonomy and privacy protocols. Restricting data retention to protect user privacy limits offline retraining and explanation auditing, while routing high-confidence alerts to human counselors increases per-query serving cost and the end-to-end latency $L_{\text{lat,total}}$, measured from the incoming message to an acknowledged alert. A local differential privacy mechanism would need a specified privacy unit, accountant, and clinical utility evaluation. The automated dispatch SLA covers producing and acknowledging an alert, not completing human counseling. Human response needs a separate deadline, staffed queue capacity, and an escalation policy for overload.

Systems lesson: Responsible AI trade-offs are hard architectural requirements, not after-the-fact policy notes. The deployment pipeline must encode stakeholder choices into explicit latency ($L_{\text{lat,total}}$) and privacy budgets before optimization begins.

What constitutes a fair outcome for one stakeholder may be perceived as inequitable by another. Similarly, decisions that prioritize accuracy or efficiency may conflict with goals such as transparency, individual autonomy, or harm reduction. These tensions are structural rather than incidental, reflecting the pluralistic societies and institutional settings in which machine learning systems operate. Fairness is a particularly prominent site of value conflict because its criteria can be incompatible (section 16.5.1.5) and choosing among them is a normative decision rather than a purely technical one, yet similar conflicts arise between interpretability and predictive performance, privacy and personalization, or short-term utility and long-term consequences.

Each deployment tier gives these conflicts a different form, adding to the tier comparison of section 16.4.1 the question of who decides. A voice assistant that personalizes on a mobile device keeps raw data local, which reduces transmission exposure without by itself establishing a privacy guarantee. Users with underrepresented usage patterns may still receive less accurate predictions, and detecting that disparity requires decisions about consent, aggregation, and access, so privacy, fairness, and personalization must be weighed against one another. A cloud credit-scoring or recommendation platform faces demands for explanations of consequential decisions while relying on complex ensembles or proprietary training data whose disclosure may be commercially sensitive or technically infeasible, which sets institutional accountability against business confidentiality. An edge camera or drone that compresses its model to meet latency and energy budgets trades efficiency against robustness. The effect must be measured on the deployed candidates under the expected shifts and threat model rather than assumed, and a more resilient candidate may exceed the memory or power budget. A TinyML device without persistent connectivity cannot absorb new fairness requirements through local retraining. Whether servicing or intermittent synchronization can deliver a revised model, and whether explanations and appeals can run off-device on preserved evidence, therefore becomes part of the value decision. In every tier the conflict extends beyond what the model optimizes to what the system can support after deployment.

The recurring systems constraint is normative pluralism, not an abstract philosophical challenge. Technical approaches such as multi-objective optimization, constrained training, and fairness-aware evaluation can help surface and formalize trade-offs, but they do not eliminate the need for judgment. Decisions about whose values to represent, which harms to mitigate, and how to balance competing objectives cannot be made algorithmically. They require deliberation, stakeholder input, and governance structures that extend beyond the model itself. Machine learning systems are not neutral tools; they embed value judgments whether those judgments are specified or assumed.

Participatory methods and the Value-Sensitive Design process introduced with value alignment (section 16.8.3) offer paths forward. Rather than treating values as parameters to be optimized after deployment, these approaches seek to engage stakeholders during the requirements phase, define ethical trade-offs explicitly, and trace how they are instantiated in system architecture. Some requirements can be satisfied together; others require a documented choice or a decision not to deploy. Making these judgments explicit gives affected stakeholders something concrete to question and revise.

Value-Sensitive Design also addresses the autonomy of the people a system serves. Applied to human control, it explicitly considers factors such as capability, complexity, misrepresentation, and the fluidity of user control. The METUX model (Motivation, Engagement, and Thriving in the User Experience) extends this thinking to six spheres of technology experience, namely adoption, interface, tasks, behavior, life, and society, each of which can support or undermine human flourishing (Peters et al. 2018). Both draw on self-determination theory's account of autonomy as the ability to act in accordance with one's own values and goals, rather than as control in a technical sense (Ryan and Deci 2000). In an ML system those requirements become concrete artifacts, namely interface affordances, user controls, feedback channels, escalation paths, and audit evidence. A recommender that optimizes engagement, for instance, can shape user preferences in ways the user cannot see. Evaluating it across the outer spheres of behavior, life, and society exposes effects that a short-term

engagement metric never measures, so preserving autonomy requires controls and evidence beyond that metric.

Feedback loops, human-AI collaboration, and value conflicts show why the technical foundations from section 16.5 through section 16.8 cannot ensure responsible AI on their own. Without organizational authority, operational telemetry, and escalation paths, a fairness metric with no owner, an explanation with no appeal process, or a monitoring signal with no remediation budget becomes evidence without action.

Checkpoint 16.3: When correct code fails in deployment

Mathematically sound, fair, and explainable models can still fail in human institutions through feedback loops that reshape the data, automation bias and institutional deference, and normative pluralism among stakeholders.

Diagnosing the human-system interaction

- ☐ **Alert adoption failure**: propose two workflow failures that could explain ignored sepsis alerts despite high offline accuracy, and identify a measurement that distinguishes them. Propose an interface or workflow change for one explanation, then state how timely action, actual error correction, workload, and patient outcomes would be evaluated.
- ☐ **Automation bias paradox**: explain why overall accuracy and override frequency do not establish the rate of actual errors corrected, and identify what evidence would justify an interface change.

Reasoning about loops and stakeholders

- ☐ **Feedback loops**: trace how a feedback loop makes a model's original prediction self-fulfilling in a predictive-policing or recommender system, and explain why a static test set fails to detect it.
- ☐ **Normative pluralism**: justify why no algorithm resolves conflicts among medical efficacy, patient autonomy, privacy, and legal compliance in the adolescent mental-health crisis-detection pipeline, and identify what the engineering team must do before optimization begins.

16.10 Governing Responsible AI at Fleet Scale

The data science team wants to hold back deployment for a month to conduct rigorous fairness audits on a new generative model. The executive team, watching a competitor launch a similar feature, demands the model be deployed by Friday. Even when engineers know how to evaluate the relevant risks, the organizational structure, budget, and business priorities must give them the time and authority to act on that knowledge.

The scenario exposes the implementation gap between technical capability and operational authority. Responsible AI methods provide necessary tools, but their effectiveness depends on organizational structures, data infrastructure, evaluation processes, and sustained commitment that extends far beyond algorithm development. Systems maintain responsible behavior over time only when those implementation supports survive deadline pressure, product incentives, and postdeployment drift.

The practical barriers to embedding responsible AI in production ML systems span the classical People-Process-Technology framework. People challenges, the subject of section 16.10.1, encompass the organizational structures, authority, role definitions, and incentives that determine whether ethical commitments translate into durable operational behavior. Technology challenges begin with the data itself (section 16.10.2), where quality gaps and acquisition costs decide what a team can fix. Process challenges cover how competing objectives are resolved (section 16.10.3) and how responsibility constraints survive scaling, maintenance, and continuous delivery (section 16.10.4). A decision framework (section 16.10.5) then turns the three dimensions into deployment-specific priorities.

These dimensions matter because responsible AI fails when organizational incentives, development workflows, and production infrastructure pull in conflicting directions. The framework connects each technical claim to the process and owner needed to act on it. A formal privacy guarantee, for example, does not implement deletion; that requires lineage and a removal workflow for the relevant data and model state. Effective implementation therefore demands systems-level strategies that embed responsibility into the architecture, infrastructure, and operating practices of machine learning deployment across all three dimensions simultaneously.

16.10.1 Organizational structures and incentives

When a modeling team is evaluated on quarterly engagement lift, data pipelines belong to a separate platform organization, user appeals land in customer support, and compliance reviews happen in legal, no single owner has both the budget and the stop-ship authority to repair a subgroup error gap or broken explanation path. The gap can then sit unfixed for months because organizational structures fragment responsibility across teams whose incentives never align on remediation. Transparency and contestability fail when responsibility is distributed without corresponding authority; within companies, research labs, and public institutions, governance must translate ethical commitments into concrete roles, escalation paths, review capacity, and authority to restrict deployment or provide remedy.

Responsible AI requires sustained investment in practices such as subgroup performance evaluation, explainability analysis, adversarial robustness testing, and the integration of privacy-preserving techniques like differential privacy or federated training. These activities can be time-consuming and resource-intensive, yet they often fall outside the formal performance metrics used to evaluate team productivity. For example, teams may be incentivized to ship features quickly or meet performance benchmarks, even when doing so undermines fairness or overlooks potential harms. When ethical diligence is treated as a discretionary task, instead of being an integrated component of the system lifecycle, it becomes vulnerable to deprioritization under deadline pressure or organizational churn.

Incentive pressure compounds when no cross-functional protocol connects detection to root-cause investigation across those team boundaries. Even when developers, users, or external auditors document disparate impact in predictions or degraded explanation fidelity, the finding can stall between model, data, user-experience, and compliance owners if no workflow assigns who must fund and execute the fix. Concerns then go unaddressed not because of malicious intent, but because the organization lacks a single escalation path with authority to coordinate mitigation.

Establishing effective organizational structures for responsible AI requires more than policy declarations. It demands operational mechanisms: designated roles with responsibility for ethical oversight, documented escalation pathways, accountability for postdeployment monitoring, and incentives that reward teams for ethical foresight and system maintainability. Some organizations also embed domain experts or user advocates in product teams to anticipate downstream impacts and evaluate value trade-offs in context.

Technical infrastructure such as versioned model registries, model cards, and audit logs must be coupled with organizational structures such as ethics review boards, model risk committees, and red-teaming[23] procedures. Their effectiveness depends on whether findings reach an authorized owner, receive a timely response, and lead to a verified correction when one is required. The existence of a committee or audit log alone does not establish that concerns have been addressed.

Operational rigor comes with a measurable cost. Public responsible-AI standards and internal review processes illustrate that governance can add release latency through impact assessments, model reviews, red-team exercises, and documentation gates. Review capacity therefore becomes part of deployment planning. If every high-risk launch needs specialized review, the review queue becomes a bottleneck just like security review or capacity approval. That cost must be budgeted alongside remediation and residual risk; preventive review can reduce exposure, but does not guarantee that every incident will be avoided or that remediation will always cost more.

Third-party and foundation-model dependencies do not transfer accountability to the supplier. A deploying organization may inherit weights, safety evaluations, data documentation, or contractual controls, but it still chooses the use, interface, retrieved context, user population, monitoring, and response process. Procurement therefore needs the same rigor as an internal model gate: documented intended and prohibited uses, evaluation scope, version and change notification,

[23] **Red teaming**: From Cold War military simulations where the "Red Team" acted as the Soviet adversary to probe US defenses. In Responsible AI, red teaming is the adversarial audit phase: specialized teams (hackers, linguists, ethicists) deliberately probe models for jailbreaks, bias, or toxic outputs before deployment. This discovery process identifies the long-tail risks that standard unit tests cannot catch.

incident disclosure, audit access, data handling, service dependencies, fallback behavior, and an exit plan. When supplier evidence is unavailable, the assurance case must record the gap as residual risk rather than silently treating a vendor assertion as a verified control.

The organization must provide these functions across the actual deployment lifecycle. A cloud team may operate continuous review and monitoring services; an embedded-device operator may rely on certification, firmware inventories, servicing, and recall. Domain experts, affected parties, and regulators remain part of the evidence and authority structure in either case. The operating arrangement must still support correction of an individual decision.

Accountability becomes visible at the point of remedy. Affected people need a discoverable channel to challenge an outcome, a record of the evidence considered, a reviewer with authority to correct the decision, a response deadline, and a way to propagate the correction into downstream systems. Aggregate model improvement is not a substitute for individual remedy, because retraining next quarter does not correct a benefit denied today. Incident procedures must therefore connect technical containment and rollback to notification, correction, compensation where applicable, root-cause analysis, and verification that the remedy reached every dependent decision record. Retirement does not end these duties: shutting down a model closes its execution path, but an appeal owner remains responsible for unresolved decisions after inference stops (section 16.2). England's 2020 exam-grading algorithm showed how quickly an aggregate method can outrun the remedy path its decisions require.

War Story 16.2: The algorithmic grading failure (2020)

Context: In August 2020, following exam cancellations during COVID-19, England's qualifications regulator (Ofqual) deployed a statistical standardization algorithm to grade ~280,000 A-level students (Ofqual 2020; Department for Education and Ofqual 2020).

Mechanism: To control grade inflation, the standardization method combined historical results with the cohort's prior attainment to estimate grade distributions, then allocated grades using teachers' student rankings. Small subject cohorts received greater weight on teacher-assessed grades. The method therefore used individual ranking evidence, but constrained its interpretation through cohort-level predictions that could differ from teachers' assessments of individual achievement.

Impact: The calculated grade was lower than the teacher-assessed grade in 39.1 percent of cases (Ofqual 2020). Public outcry and student protests prompted a reversal of the grading policy days later.

Fix: Students were awarded the higher of the center-assessed grade and the calculated grade, replacing the earlier standardization policy.

Systems lesson: An aggregate prediction does not establish that every individual outcome is justified. An appeal channel must specify which evidence can change the decision, who can review it, and how quickly a correction can take effect; the existence of a procedural appeal alone does not answer those questions.

The responsibility for ethical system behavior is distributed across multiple constituencies, including industry, academia, civil society, and government. The resulting governance hierarchy surrounds the AI system with four concentric tiers (figure 16.14). The development team executes technical controls (fairness testing, model cards, red-teaming). The organization governs operational processes (impact assessments, ethical review boards). Industry establishes consensus baselines (ISO standards, external audits). Government and society enforce legal boundaries (GDPR, the EU AI Act, sector-specific statutes) (Shneiderman 2022, 2020). Within organizations, this distribution must be mirrored by mechanisms that connect technical design with strategic oversight and operational control. Without these linkages, responsibility becomes diffuse, and well-intentioned efforts may be undermined by systemic misalignment.

Technical excellence and regulatory compliance are therefore insufficient on their own. Responsible AI requires aligning ethical objectives with the institutional structures through which machine learning systems are designed, deployed, and maintained, so that responsibility is embedded in the

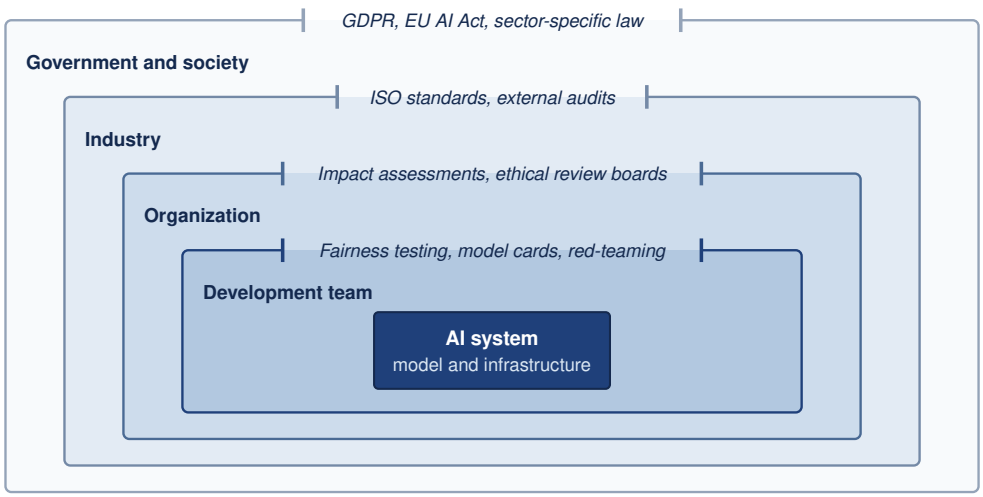

Figure 16.14: Layers of Responsibility: Effective human-centered AI implementation requires shared accountability across nested layers. From the core Development Team outward through the Organization, Industry, and Government/Society, each surrounding layer enforces norms the inner layers cannot self-impose. External anchors, including ISO AI management standards, fairness testing, model cards, red-teaming, impact assessments, GDPR, the EU AI Act, and sector-specific law, bind these layers to practice (International Organization for Standardization and International Electrotechnical Commission 2023b; European Parliament and Council of the European Union 2016; European Parliament and Council of the European Union 2024).

organization that governs deployment as well as in the model itself. The same ownership problem becomes concrete at the data layer, where teams can know how to audit a dataset and still lack the authority, access, or infrastructure needed to change it.

16.10.2 Data constraints and quality gaps

A clinical imaging model trained on a handful of urban teaching hospitals can post strong headline accuracy while failing on rural and elderly patients whose scanners, protocols, and demographics were missing from the training corpus, yet acquiring and harmonizing replacement images across dozens of partner hospitals requires years of legal, institutional, and schema negotiations. Improving data pipelines remains one of the hardest implementation challenges in practice, even when developers and researchers understand the need for representative coverage, accurate labels, and historical-bias mitigation. Responsibility for data is often split across teams, locked inside legacy systems, or governed by external institutions that a modeling team cannot unilaterally change.

Data engineering principles, including data validation, schema management, versioning, lineage tracking, and quality monitoring, provide the technical foundation for addressing these challenges. However, applying these principles to responsible AI introduces additional complexity. Fairness requires assessing representativeness across demographic groups, bias mitigation demands understanding historical data collection practices, and privacy preservation constrains which validation techniques are permissible. The organizational challenges described here reflect the gap between having robust data engineering infrastructure and using it effectively to support responsible AI objectives.

Subgroup imbalance, label ambiguity, and distribution shift, each of which affect generalization and performance across domains, are well-established concerns in responsible ML. These issues often manifest in the form of poor calibration, out-of-distribution failures, or demographic disparities in evaluation metrics. However, addressing them in real-world settings requires more than technical knowledge. It requires access to relevant data, institutional support for remediation, and sufficient time and resources to iterate on the dataset itself. Teams may lack the authority or infrastructure to modify a training dataset even after discovering performance disparities. Versioning and lineage make controlled correction possible by identifying the affected datasets and downstream models; they do not remove the need to acquire replacement data, retrain, and revalidate those models.

In domains like healthcare, education, and social services, these barriers are especially pronounced because data acquisition is governed by legal constraints, privacy regulations, and cross-organizational contracts. Closing the rural-hospital coverage gap in the opening clinical scenario requires negotiating data-sharing agreements with external partners, aligning on feature standards, and resolving inconsistent labeling practices across institutions, and the logistical and financial costs can be prohibitive even when all parties agree on the need for improvement (notebook 16.8).

Efforts to collect more representative data may also run into ethical and political concerns. In some cases, additional data collection could expose marginalized populations to new risks. This paradox of exposure, in which the individuals most harmed by exclusion are also those most vulnerable to misuse, complicates efforts to improve fairness through dataset expansion. For example, gathering more data on nonbinary individuals to support fairness in gender-sensitive applications may improve model coverage, but it also raises serious concerns around consent, identifiability, and downstream use. Teams must navigate these tensions carefully, often without clear institutional guidance.

Napkin Math 16.8: The representation cost

Problem: A medical imaging model trained on data from 5 hospitals achieves 94 percent accuracy overall but only 78 percent on underrepresented populations (rural patients, elderly patients, patients with darker skin tones). A proposed acquisition plan seeks data from 50+ hospitals across diverse geographies, demographics, and equipment types. This budget is an illustrative proposal, not a sample-size requirement derived from the accuracy gap; validation must establish whether it improves the relevant outcomes.

Math:

1. **Data acquisition cost**: $50–$200 per labeled medical image, with 100,000 images budgeted per underrepresented subgroup.
2. **Representation-cost scenario**: For 10 subgroups, at the midpoint of the cost range, 10 times 100,000 images per subgroup times $125 per image yields $125M in data acquisition alone. This calculation assumes distinct purchased examples for each allocation; overlapping demographic groups require deduplication or explicit shared-cost allocation.

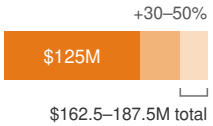

Harmonizing scanners, protocols, and labels adds 30–50 percent to the acquisition cost.

3. **Data harmonization**: Data harmonization across different scanners, protocols, and labeling conventions adds 30 percent–50 percent overhead, bringing the total to $162.5M–$187.5M.

Systems insight: The acquisition budget makes coverage goals explicit, but spending does not guarantee that an accuracy gap closes. Teams should prioritize identified coverage deficiencies using harm, data quality, privacy, and expected benefit criteria, then validate the resulting model on the intended populations.

Upstream biases in data collection systems can persist unchecked even when data is plentiful. Many organizations rely on third-party data vendors, external APIs, or operational databases that were not designed with fairness or interpretability in mind. For instance, Electronic Health Records, which are commonly used in clinical machine learning, often reflect systemic disparities in care, as well as documentation habits that encode racial or socioeconomic bias (Himmelstein et al. 2022). Teams working downstream may have little visibility into how these records were created, and few levers for addressing embedded harms.

Improving dataset quality is often not the responsibility of any one team. Data pipelines may be maintained by infrastructure or analytics groups that operate independently of the ML engineering or model evaluation teams. This organizational fragmentation makes it difficult to coordinate data audits, track provenance, or implement feedback loops that connect model behavior to underlying data issues. In practice, responsibility for dataset quality tends to fall through the cracks, recognized as important, but rarely prioritized or resourced.

Addressing these challenges requires long-term investment in infrastructure, workflows, and cross-functional communication. Data validation and automated audits can identify defects, while datasheets and the Data Nutrition Project[24] document dataset properties. Model cards connect model evaluation and intended use to that evidence. These tools help only when teams have the mandate and support to act on their findings.

Improving data quality is therefore a question of how responsibility for data is assigned, shared, and sustained across the system lifecycle, not merely a matter of better tooling. Once that responsibility is in place, teams still have to decide which responsible-AI objective should dominate when valid goals conflict.

[24] **Data Nutrition Project**: A nonprofit initiative that develops standardized dataset documentation modeled on nutrition labels, summarizing provenance, composition, distribution, and known issues for datasets used in ML pipelines. The project's labels make dataset risks legible to downstream engineers who lack access to the original collection process, which bridges the organizational fragmentation described here.

16.10.3 Balancing competing objectives

When a modeling team tunes for aggregate accuracy, a compliance team injects differential-privacy noise that degrades minority-group recall, and a serving team prunes the network to hit a 15 ms latency target, their uncoordinated local optimizations can break the deployment's fairness and explanation guarantees. Every team nonetheless met its own metric. Machine learning system design is often framed as optimizing a single loss or accuracy objective, yet production systems must balance accuracy against fairness, interpretability, robustness, privacy, and resource efficiency, objectives that routinely pull in opposite directions. The accuracy-interpretability, personalization-fairness, and privacy-utility conflicts of section 16.4.1 recur here, and each must be measured on the path where its mechanism actually executes rather than assumed. Those trade-offs are not purely technical; they reflect deeper normative judgments about what a system is designed to achieve and for whom, as explored in detail in section 16.9.4.3. Responsible ML development requires making these judgments explicit, evaluating them in context, and subjecting them to stakeholder input and institutional oversight.

What makes this challenge particularly difficult in implementation is that these competing objectives are rarely owned by a single team or function. Performance may be optimized by the modeling team, fairness monitored by a responsible AI group, and privacy handled by legal or compliance departments. Without deliberate coordination, system-level trade-offs can be made implicitly, piecemeal, or without visibility into long-term consequences. Over time, the result may be a model that appears well-behaved in isolation but fails to meet its ethical goals when embedded in production infrastructure.

Balancing competing objectives requires transparency, deliberation, and cross-team alignment alongside technical fluency. Systems must be designed to surface trade-offs rather than obscure them, to make room for constraint-aware development rather than pursue narrow optimization. In practice, this may require redefining what "success" looks like, not as performance on a single metric, but as sustained alignment between system behavior and its intended role in a broader social or operational context.

Organizational structures, data quality, and competing objectives explain why technical competence alone does not secure responsible behavior. A team can understand fairness metrics and still lack authority to change a dataset or delay a release. When modeling teams optimize performance, compliance teams address privacy, and product teams prioritize engagement independently, their local decisions may violate a shared requirement. Clear ownership across those boundaries connects evaluation to action, while data infrastructure and deliberative review make the evidence and trade-offs inspectable. Maintaining those connections becomes harder as deployments scale and their data, models, and teams change.

16.10.4 Scalability, maintenance, and evaluation standards

Safeguards built for a release review frequently degrade or disappear as a prototype scales into continuous production: fairness checks run on a static dataset never join the live ingestion pipeline, explanation methods evaluated during model selection never reach a user-facing interface, and privacy constraints enforced during initial training are overlooked in postdeployment monitoring or scheduled updates. Closing the gap between what works in a one-time evaluation and what survives continuous delivery is a core implementation challenge for responsible AI.

Production environments introduce new pressures that reshape system priorities. Models must operate across diverse hardware configurations, interface with evolving APIs, serve millions of users

with low latency, and maintain availability under operational stress. These constraints demand continuous adaptation and rapid iteration, often deprioritizing activities that are difficult to automate or measure. Responsible AI practices, especially those that involve human review, stakeholder consultation, or post-hoc evaluation, may not be easily incorporated into fast-paced DevOps[25] pipelines.

Maintenance introduces further complexity. Machine learning systems are rarely static. New data is ingested, retraining is performed, features are deprecated or added, and usage patterns shift over time. In the absence of rigorous version control, changelogs, and impact assessments, it can be difficult to trace how system behavior evolves or whether responsibility-related properties such as fairness or robustness are being preserved. Organizational turnover and team restructuring can erode institutional memory. Teams responsible for maintaining a deployed model may not be the ones who originally developed or audited it, leading to unintentional misalignment between system goals and deployed behavior. These issues are especially acute in continual or streaming learning. Drift calls for investigation and a response matched to its consequences, which may mean restricting use or invoking fallback while a revised model undergoes validation, rather than immediately updating production weights.

These challenges are magnified in multi-model systems and cross-platform deployments. A recommendation engine may consist of dozens of interacting models, each optimized for a different subtask or user segment. A voice assistant deployed across mobile and edge environments may maintain different versions of the same model, tuned to local hardware constraints. Coordinating updates, ensuring consistency, and sustaining responsible behavior in such distributed systems requires infrastructure that tracks governing values and constraints alongside code and data. Addressing these challenges means treating responsible AI as a lifecycle property rather than a one-time evaluation: audit hooks, metadata tracking, and monitoring protocols built into infrastructure, with documentation and accountability structures that survive team handoffs.

Maintenance also exposes a gap in how responsible-AI evidence is standardized. A model evaluated under demographic parity on one benchmark has not thereby satisfied equalized odds in another deployment. The criterion, population, labels, threshold policy, and uncertainty belong to the claim. Changing any of them can invalidate the evidence even when the weights remain unchanged. Standardization must make those conditions explicit so teams can compare results and identify what needs revalidation.

The evaluation boundary must also match the deployment boundary. A model-only result can miss effects introduced by feature transformations, request routing, caching, or human review. Separate fairness, explanation, and compliance tools therefore need compatible identifiers and records linking their results to the same model, data, policy, and workflow versions. Otherwise, an assurance case can combine individually valid results that describe different systems.

Reusing that evidence after launch requires lifecycle controls. Production observations must be comparable with the acceptance criteria, and material changes must trigger a documented revalidation scope. Evidence must therefore remain interpretable across teams, components, and time. The implementation decision is then which obligations take priority for the intended use, which controls can satisfy them, and which authority may accept the remaining risk.

16.10.5 Implementation decision framework

An intensive-care diagnostic model, a 10,000-query-per-second consumer recommender, and a milliwatt microcontroller sensor cannot run the same safeguards on the same path. The ICU model needs clinician-legible explanations and contested-case routing. The recommender needs continuous slice-drift monitoring and feedback-loop breakers. The microcontroller must lock in its safety envelope and energy budget before firmware ships. Practitioners therefore need a systematic way to prioritize responsible AI controls by deployment context and stakeholder risk.

Four decision heuristics guide these trade-offs in practice. Where multiple principles conflict, engineering teams must engage domain stakeholders to establish which failure modes present the most severe harms. Under tight computational budgets, teams must first identify mandatory requirements, then allocate evaluation and response capacity according to risk. Sampling is adequate only if it can detect the relevant failures, and expensive audits need not run on the inference path. Whenever the deployment context changes, as when a cloud model migrates to edge execution

25 **MLOps (machine learning operations)**: ML CI/CD pipelines must handle data versioning, training reproducibility, and A/B testing of algorithm changes beyond traditional software concerns. High-velocity ML organizations may deploy model or configuration changes far more often than responsible-AI reviews can be completed manually, creating a velocity gap in which deployment cycles measured in hours compete against ethical validation requiring days or weeks. This tension explains why responsible AI commitments present at the prototype stage can be deprioritized as systems scale.

with intermittent centralized monitoring, teams must re-evaluate principle priorities, establishing local fallback safeguards and evidence-synchronization plans for constrained connectivity. Finally, when stakeholder values diverge, engineers must document trade-offs explicitly and construct contestability mechanisms that permit affected individuals to challenge automated outcomes. As the adolescent mental health crisis-detection example in section 16.9.4.3 demonstrated, such value tensions require deliberate institutional policy rather than mathematical optimization alone.

Each deployment context suggests candidate priorities and trade-offs for deployment-specific review (table 16.10). Applicable obligations, the threat model, and accountable stakeholder decisions determine which controls and trade-offs are acceptable.

Table 16.10: Practitioner Decision Framework: Prioritizing responsible AI principles based on deployment context, showing candidate principles, controls, and trade-offs for different system types. Deployment-specific review must establish applicable requirements and justify the selected controls; no row waives mandatory obligations.

Deployment context	Primary principles	Implementation priority	Trade-offs to evaluate
High-Stakes Individual Decisions (healthcare diagnosis, credit and loans, criminal justice, employment)	Fairness, explainability, accountability	Mandatory fairness metrics across protected groups; explanations for negative outcomes; human oversight for edge cases	A measured accuracy budget for interpretability; bounded explanation latency; higher computational cost
Safety-Critical Systems (autonomous vehicles, medical devices, industrial control)	Safety, robustness, accountability	Threat-model-specific evaluation; bounded certification where applicable; fail-safe mechanisms; comprehensive logging	Significant validation and adversarial-training overhead; conservative confidence thresholds; redundant inference
Privacy-Sensitive Applications (health records, financial data, personal communications)	Privacy, security, transparency	Data minimization and appropriate consent mechanisms; evaluate local processing and differential privacy with stated (ϵ, δ), privacy unit, and accountant	Privacy-utility trade-offs; higher client-side compute; limited model updates; reduced personalization
Large-Scale Consumer Systems (content recommendation, search, advertising)	Fairness, transparency, safety	Bias monitoring across demographics; explanation mechanisms; content-policy enforcement; feedback-loop detection	Sampled or asynchronous explanations at scale; serving-path latency budgeted for fairness checks; investment in monitoring infrastructure
Resource-Constrained Deployments (mobile, edge, TinyML)	Privacy, efficiency, safety	Local inference; data locality; input validation; graceful degradation	Expensive audits moved off-device, while local safety checks, fleet or servicing audits, update authority, and retirement criteria are retained
Research and Exploratory Systems (internal tools, prototypes, A/B tests)	Transparency, safety (harm prevention)	Documented known limitations; restricted user populations; monitoring for unintended harms	Evaluation depth matched to potential harm and applicable obligations; observability and bounded iteration

Lighthouse 16.1: Lighthouse archetypes: Evidence and decision authority

The lighthouse workloads in section 1.6.1 require different evidence because they expose different harms, and they assign decision authority to different operating paths.

Workload	Minimum decision evidence	Required authority
Archetype A (GPT-4/Llama-3)	Capability- and use-specific safety evaluations, harmful-output and misuse tests, uncertainty and fallback behavior, and versioned deployment evidence	Model and safety owners approve the allowed use envelope; service owners can halt or roll back when monitoring invalidates it
Archetype B (DLRM at Scale)	Disaggregated ranking quality, exposure and error rates, feedback-loop analysis, and online effects under the serving policy	Product and responsible-AI owners choose the welfare and fairness criteria; operators can disable a model or policy that breaches them
Archetype C (Federated MobileNet)	Per-device and subgroup quality, participation bias, privacy accounting, update integrity, and behavior across client versions	Device, privacy, and model owners approve participation and update policy; fleet control can suspend aggregation or revert clients

A metric becomes a control only when the named authority can act on it. Evidence that is sufficient for a recommender policy review does not establish safety for a general-purpose language model, and aggregate federated accuracy does not authorize an update whose participation or privacy assumptions have failed.

The framework provides starting guidance, but it should be read as a triage mechanism rather than a complete governance program. Legal obligations depend on intended use, jurisdiction, and organizational role, not whether inference runs in a cloud or on a microcontroller (European Parliament and Council of the European Union 2024; European Parliament and Council of the European Union 2016). The framework identifies which constraints deserve first attention, with reassessment as deployment conditions change.

The lifecycle assigns human authority over deployment even when individual actions are automated. Engineers investigate bias, operators review fairness metrics, and developers respond to drift. Some systems must also act faster than humans can review an individual decision. Autonomous vehicles respond in milliseconds; trading algorithms execute thousands of transactions before human review is possible; content moderation systems process billions of posts daily. For these systems, the implementation framework must be paired with the value-alignment and fallback controls of section 16.8.3, which keep objectives aligned with human values even when operation runs beyond continuous human supervision.

Systems Perspective 16.2: Responsibility is infrastructure, not a feature

Retrofitting responsible AI controls can require architectural redesign and revalidation, just as adding fault tolerance can. The structural costs cataloged in table 16.5 each carry a provisioning consequence. Bias-drift monitoring needs a sampling and analysis budget. Expensive or exact explanation workloads may need a dedicated asynchronous worker fleet, whereas model-specific TreeSHAP, sampling, caching, or selective generation may fit smaller budgets. DP-SGD requires budgeting its clipping, noise, and accounting work under the chosen implementation. Impact assessments need review capacity in the release process. These are not *optional add-ons* but *load-bearing components* of production ML systems. In high-risk or regulated deployments, failing to provision the evidence required by the assurance case can block release or force redesign even when core model performance is strong.

16.11 Fallacies and Pitfalls

Responsible AI involves counterintuitive trade-offs where ethical principles conflict mathematically and technically. Practitioners from traditional software backgrounds often assume ethical guidelines translate directly to implementation without recognizing the impossibility theorems and computational costs involved. These fallacies and pitfalls are misconceptions that hide the structural costs of responsibility or mistake a local technical result for system-level assurance. They lead to deployed systems that appear fair in development but violate fairness criteria in production or impose prohibitive computational overhead.

Fallacy: *Bias will disappear with more data and better algorithms.*

More representative data and better algorithms can reduce some disparities, but they do not determine whether the prediction target is appropriate. The healthcare algorithm described in section 16.5.1 underestimated Black patients' health needs at a given risk score because expenditure was an imperfect proxy for need. Improving prediction of that proxy would not by itself repair the target. Separately, the score-balance impossibility result of section 16.5.1.5 shows that, when base rates differ and prediction is imperfect, no score can be calibrated within each group and also balanced across groups in both outcome classes, a limit that no data or algorithm improvement removes. Engineers must distinguish a remediable data or modeling defect from incompatible requirements, then use stakeholder deliberation to justify the criteria the deployment will pursue.

Pitfall: *Treating explainability as an optional feature to be added after core functionality works.*

This approach ignores the displacement of overhead (principle 14). Deferred explanation work does not vanish; it moves to the request path, the memory budget, or a review queue, because producing attribution evidence about model behavior demands compute, storage, and interfaces that constrain the architecture. Approximate SHAP explanations are not compute-free (table 16.5); even a bounded explanation budget can increase request latency and memory pressure. In this sizing example, a recommendation system serving 100 ms latency requirements at 10,000 QPS cannot retrofit explanation work without revisiting the SLA. The illustrative explanation budget adds 50 ms to 200 ms per request, increasing total latency to 150 ms–300 ms. This requires 1.5–3× capacity only under the simplifying assumption that occupied resource time grows in direct proportion to latency, at fixed utilization and parallelism. The assurance case must specify the audience, claim, method, fidelity test, and delivery budget, then choose an architecture such as model-specific TreeSHAP, sampled or cached attribution, selective asynchronous generation, or an inherently interpretable model. Additional replicas can preserve throughput under that assumption, but they do not restore the original per-request latency if explanation remains on the serial path. None of these attribution paths exposes a model's internal reasoning or establishes causation.

Fallacy: *Achieving one fairness metric guarantees overall system fairness.*

Practitioners often optimize for demographic parity, which equalizes approval rates across groups, and then assume fair treatment follows. In reality, fairness metrics can conflict mathematically. In the credit-scoring scenario of notebook 16.4, reaching approval parity cost 4 percentage points of accuracy and left approvals in the two groups carrying different default risk. The score-balance impossibility result shows that, with unequal base rates and imperfect prediction, groupwise calibration and equal mean scores across groups in both outcome classes cannot all hold; demographic parity can add a separate selection-rate constraint. Passing one metric therefore leaves other criteria to evaluate. A metric disparity is evidence for investigation, while a legal conclusion depends on the applicable rule and the facts of the deployment.

Pitfall: *Implementing fairness constraints without analyzing threshold trade-offs and calibration impacts.*

Moving a threshold can equalize one metric while leaving the underlying score calibration unchanged. In production, the threshold is an operational policy embedded in serving code, human review, audit records, and downstream decisions rather than a standalone cutoff in a notebook. Group-specific thresholds can satisfy equal true positive rates, but they also change decision semantics. If the Group A threshold is 0.75 and the Group B threshold is 0.45, as in figure 16.5, identical scores can yield different decisions, and the positive predictive value of approved applicants can diverge across groups even when the underlying score model remains calibrated. A loan officer seeing an approval cannot know whether it represents the same repayment-risk meaning across groups without documentation of the group-specific threshold policy. Production systems therefore need documented metric selection, uncertainty estimates, staff training, legal review, and audit trails explaining why identical scores can yield different decisions. No generic training objective resolves incompatible fairness definitions; the appropriate response is to choose and justify the criteria for the deployment context, quantify the resulting trade-offs, and monitor whether the approved policy remains valid.

Fallacy: *Monitoring dashboards create accountability by themselves.*

Dashboards can expose fairness drift, explanation failures, or abuse trends, but they do not decide who must act. Teams often instrument fairness, safety, or privacy metrics because a policy requires visibility, then leave the response path undefined, which creates a false sense of control. A metric that no team owns is not a control; it is a log entry. The system accumulates evidence without changing behavior, so alerts age out, fairness regressions become accepted baselines, and affected users receive no route to contest the decision. Production responsibility requires binding each monitoring signal to an owner, a threshold, an escalation path, and a remediation budget, whether that is a release gate, an incident owner, rollback authority, or human-review capacity.

Pitfall: *Treating responsible AI as pure compliance overhead.*

This perspective misses the value of identifying and responding to risk. Differential privacy can enable organizations to work with sensitive data under clearer privacy guarantees, but it also imposes privacy-utility trade-offs that must be planned in the training budget. Fairness-aware training and monitoring can support that risk-management process. A compliance screen such as the employment-selection threshold discussed in section 16.5.2.1 checks one context-specific selection-rate ratio, not a universal legal test, and it cannot see a disparity hidden in the label itself. Monitoring need-adjusted outcomes by group could reveal a disparity such as the healthcare algorithm's underestimate of Black patients' needs. Reducing the resulting harm still requires an owner with authority, capacity, and a validated remedy; detection alone does not protect affected people.

Treating a local metric or control as sufficient system assurance is the common failure behind these examples. A mathematical fairness constraint is only responsible when the system also accounts for downstream population effects, documentation requirements, operational capacity, and legal exposure.

16.12 Summary

Responsible AI connects the fleet's technical controls to justified decisions about deployment and remedy. Fairness metrics expose disparities, privacy and robustness controls support bounded claims, and explanations provide evidence that people must be able to inspect or contest. None of these mechanisms establishes responsibility on its own. A versioned assurance case connects the intended use, affected population, evidence, assumptions, and accountable authority, while monitoring and revalidation keep those claims tied to the system that actually runs.

Responsible AI is a systems engineering concern because its controls reshape the production architecture. Fairness monitoring consumes compute and can affect latency. Explainability methods such as SHAP and LIME add inference work that must be factored into capacity planning and SLO compliance. System prompts, model versions, and audit logs require ownership, versioning, and CI/CD gates. These costs must be budgeted during design rather than added after deployment.

Selective explanation generation illustrates the C^3 choice. Producing an explanation synchronously adds compute to the request path and may add communication to retrieve features; moving it to an asynchronous worker protects serving latency but adds coordination among the decision, evidence record, and appeal workflow. The bottleneck then moves from inference capacity to evidence freshness or reviewer throughput. The acceptable architecture depends on who needs the explanation and how quickly that person must be able to act on it. Read through the Fleet Law (equation 1.3), which sums per-device compute, communication, and synchronization time and subtracts the time that overlaps, private or adversarial training can add compute through mechanism-specific operations or additional evaluations. Launch approvals and revalidation gates act like a synchronization barrier at release time, adding waits that replication alone cannot remove. Asynchronous evidence generation permits overlap only when the evidence still arrives before the decision or remedy that depends on it.

Key Takeaways: Responsibility needs evidence and authority

- **Responsibility gates deployment**: In high-risk or regulated settings, accuracy and latency cannot substitute for documentation, audit, and ownership. A versioned assurance case justifies one stated use, and a material change expires its approval.
- **Fairness requires justified criteria**: The running loan model violates every criterion at once, and the disaggregated-metrics table shows a second loan model with a 0 percentage-point false-positive gap while still carrying 15 percentage-point approval and 30 percentage-point true-positive gaps. The impossibility result shows that, with unequal base rates and imperfect prediction, no score can be calibrated within each group and also give equal mean scores across groups in both outcome classes, so the chosen criteria need an explicit justification.

- **A gap is evidence, not a verdict**: A subgroup disparity supports a decision only with denominators, confidence intervals, and intersectional coverage, and even a significant gap does not identify its cause. A causal claim needs an explicit structural model and assumptions that observational data cannot supply.
- **Safeguards are budgeted capacity**: Fairness monitoring, audit logs, explanations, private training, and generative guardrails such as system prompts and output filters each consume compute, storage, or review capacity on a specific path, and each is a versioned artifact. Where a control runs decides its cost and whom it protects.
- **Deletion must reach the weights**: A privacy guarantee bounds exposure but does not remove a record from trained parameters. Removing its influence needs lineage, a removal method with evidence, and an architecture, such as sharded training, chosen before training begins.
- **Deployed models reshape their data**: Predictions change the populations, behaviors, and labels that later training sees, so a static test set cannot detect a self-fulfilling loop. Human oversight helps only when it is designed against both automation bias and alert fatigue.
- **Evidence needs an owner and a remedy**: Dashboards and explanations protect people only when an owner can investigate, roll back, or escalate, and an affected person can challenge a decision and obtain correction. Every measured obligation needs an accountable operating path.

Quantification can show which groups bear a policy's errors and what its safeguards cost, but it cannot decide whose interests should take priority. The engineering responsibility is to make those choices explicit, preserve the evidence behind them, and provide authority to revisit them when the system or its consequences change. Governance makes the fleet's objective itself accountable, as well as the machinery used to pursue it.

What's Next: From responsibility to synthesis

Responsible governance joins performance, security, robustness, and sustainability as a constraint the fleet must satisfy before it may run. The four parts have now traversed the fleet from physical silicon, networks, and storage in Part I, through distributed coordination and recovery in Part II and fleet efficiency and serving in Part III, to the governance safeguards of Part IV. Each part examined its constraints largely one at a time, yet a production fleet couples them all. An optimization at the kernel layer reverberates into cluster power budgets, and a governance rule reshapes the serving architecture. The final chapter returns to the scale moment where this book began and follows the binding constraint as it leaves the node. It condenses the four parts into six principles of distributed ML systems, read through the Fleet Law and placed on the fleet stack.

↪ System synthesis, cross-layer co-design, and emerging fleet frontiers are developed in chapter 17.

SYNTHESIS

17

Conclusion

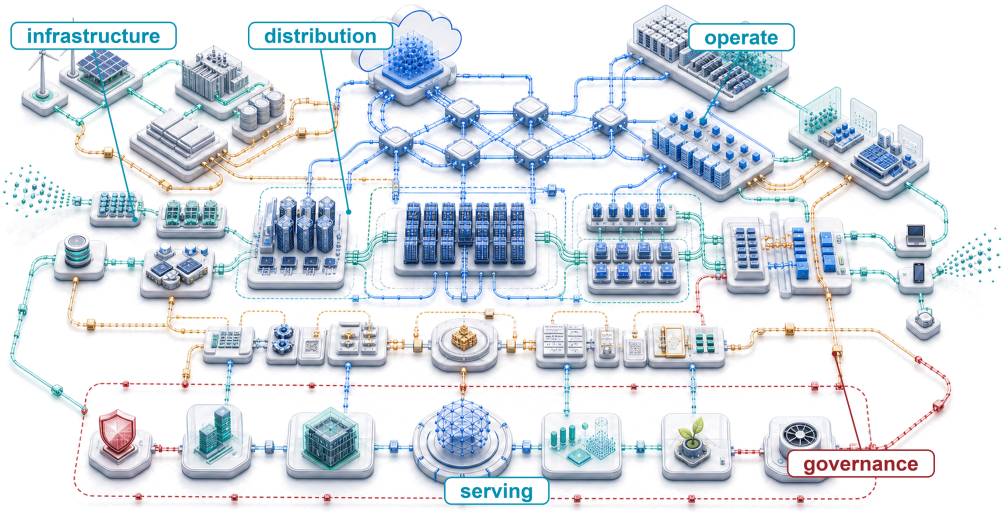

17.1 Distilling Distributed ML Systems
17.2 Six Principles of Distributed ML Systems
17.3 The Complete Production System
17.4 Competencies Mastered
17.5 Where the Binding Constraint Moves Next
17.6 Fallacies and Pitfalls
17.7 Summary

Purpose

Why must the unit of machine learning systems design become the fleet rather than the model?

Machine learning at fleet scale is judged by the whole system, and a local improvement can move cost into a layer that its designer did not measure. Faster computation can expose a network bottleneck; larger jobs can lose more work to failure; a service that meets latency targets can still lack evidence required for deployment. The model remains central, but its accuracy and execution speed do not determine whether the complete system delivers useful work. Adding accelerators is therefore a design decision about shared resources and obligations, rather than a purchase of more arithmetic. If communication and recovery consume the saved execution time, the expansion can increase cost without improving progress. If the resulting service cannot preserve availability or account for its effects on users, successful training does not establish readiness for deployment. Single-node reasoning remains necessary because every participant still has local limits, but it cannot settle decisions whose consequences cross machines, teams, and operating boundaries. The engineering responsibility is to connect those consequences before accepting the design. That requires evidence about what improves, what becomes more constrained, and who owns the costs transferred elsewhere. A defensible decision may expand the fleet, change how work is distributed, or postpone release until the operating requirements can be met. In C^3 terms, the task is to follow the binding compute, communication, or coordination constraint across layers and keep its displaced costs visible.

Learning Objectives

- Distill this book into the six principles and a constraint map for distributed ML systems
- Diagnose a fleet design's binding C^3 term using measured execution costs and failure exposure
- Analyze how infrastructure, distributed training, inference, operations, and governance interact across the fleet stack
- Translate security, fairness, privacy, carbon, and reliability obligations into measurable design constraints
- Defend a fleet expansion or model-release decision using fitted scaling-law gains, operational measurements, and governance evidence

17.1 Distilling Distributed ML Systems

A faster accelerator can finish its local work sooner yet leave the training job waiting on the same congested collective. Scale creates qualitative change because the binding constraint can leave the node. It moves into the fabric that carries gradients, into the recovery machinery that restores state after routine failure, into the service contract that decides whether capacity reaches users in time, and into the obligations that decide whether the service may run at all. The book opened on that transition, the scale moment, and on the empirical scaling laws of section 1.3 that keep forcing it. Because each further reduction in loss demands a multiplicative increase in compute, demand for fleet capacity does not settle at the size where the last design worked. Each of the four parts since has answered a share of the question that transition raised by naming a constraint, the measurement that exposes it, and the mechanism that moves it.

↩ The scale moment and the qualitative shift to fleet-scale systems are established in section 1.1.

A frontier-model production system can span thousands of accelerators, with fabric links, storage tiers, schedulers, serving replicas, security controls, and governance checks supporting the same workload. Any one layer can bind the entire system. A slow checkpoint in the storage tier can waste an entire training window. A congested fabric across the interconnect can erase scaling gains from added compute. Underprovisioned capacity in the serving pool can turn model quality into user-visible latency. An unmeasured responsibility constraint can block deployment even after the technical system works. Coupling that crosses machines, teams, and operating boundaries is what separates foundational ML engineering from distributed ML engineering. The single-node foundation, organized by the D·A·M lens of data, algorithm, and machine, already couples models, data pipelines, execution, and deployment constraints within one machine. Distributed ML engineering extends that reasoning to shared infrastructure, distributed protocols, and governance across a fleet.

↪ The single-machine foundation of data, algorithm, and machine is collected for reference in Appendix A.

The physical substrate supplies the rates that higher layers inherit. Accelerator topology, data-loading throughput, and network contention constrain how the Distribution band partitions, synchronizes, and recovers work. The Serving band converts that coordinated execution into a service with latency, availability, and cost obligations. The Governance band determines whether the service remains secure, robust, sustainable, and accountable. Six principles state these inherited constraints as engineering claims, in the order of the book's four parts.

17.2 Six Principles of Distributed ML Systems

When those inherited constraints surface in production, as a training run that stalls every few hours on a failed component or a serving pool whose tail latency spikes across a fan-out query, the failure comes from the fleet's physical and institutional structure rather than an operator mistake. The six engineering practices introduced with the fleet stack (instrument first, reserve headroom for recovery and tail latency, expect bottlenecks to migrate, treat failure as steady state, compound small efficiency gains, and co-design hardware with algorithms) describe how a fleet engineer responds to those realities, whereas the six principles of distributed ML systems describe what the fleet's physical and institutional constraints make true, whatever the practice.

The six principles condense the part-level principles stated at the opening of each part: Part I's fleet principles condense into the first, Part II's distributed ML principles into the second and third

(while its empirical scaling law supplies the demand side of the growth decision in section 17.3), Part III's fleet-in-production principles into the fourth, and Part IV's responsible fleet principles into the fifth and sixth. Where this chapter cites one of those part-level principles as "principle" followed by a number, the number refers to its box in the part opener. The six themselves are named in words, and table 17.1 pairs each with its practice.

Each principle also explains why its paired practice is necessary rather than merely prudent. A fleet that loses training progress to routine interruptions must budget recovery as part of normal operation, which is why treating failure as steady state is a practice. Because a bottleneck migrates whenever one term shrinks, an optimization judged only by the term it shrinks can leave the step no faster, which is why the engineer must expect bottlenecks to migrate and measure the layers they move into.

↩ The six engineering practices, from instrument first to co-design hardware with algorithms, are established with the fleet stack in section 1.6.

The Fleet Law, the distributed counterpart of the single-machine iron law, connects these principles to execution costs. It is the step-time equation of equation 1.3, unchanged:

$$T_{\text{step}}(N) = \frac{T_{\text{compute}}}{N} + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}} \tag{17.1}$$

Infrastructure sets the per-device rates behind T_{compute}/N and the link bandwidths behind $T_{\text{comm}}(N)$. In C³ terms, the first term is Compute, $T_{\text{comm}}(N)$ is Communication, the time spent moving bytes, and $T_{\text{sync}}(N)$ is Coordination, the time spent waiting at barriers and on stragglers, plus each step's share of run-level costs such as scheduling and recovery. Neither of the last two is guaranteed to shrink as N rises. Failure lands in the same terms. Checkpoint pauses, detection timeouts, and restarts add to $T_{\text{sync}}(N)$ averaged over the run. Asynchronous checkpointing and warm restart move part of that cost into T_{overlap}, the time hidden behind useful compute, and a rollback adds the lost steps back to the work the run must still complete. Serving applies the same accounting to each request, where queueing and routing delay form the coordination share of its latency. Responsibility and sustainability add no new term, but they spend the existing ones. Verification work occupies compute, approval gates add coordination that no replica count removes, and every second of step time costs joules at the fleet's average power. Meeting those obligations also requires evidence that execution time cannot supply, such as subgroup error rates, privacy accounting, and energy measured over a stated boundary.

↪ The term-by-term decomposition of the Fleet Law is developed in section B.3.1.

Placed on the fleet stack, the six principles follow its four bands (figure 17.1). The Infrastructure band determines physical capability. The Distribution band partitions execution and therefore carries both communication and failure recovery, one principle per row. The Serving band bounds runtime delivery by its contract. The uppermost Governance band has two rows. Its lower row, Assurance, holds security and robustness, and its top row is split into a governance half and a sustainability half. The band's two principles divide it by obligation rather than strictly by row. Responsibility covers every obligation owed to the people the system affects, so it spans the whole Assurance row and the governance half of the top row, where chapter 16 develops fairness and accountability. Sustainability covers the resource obligation alone, in the other half of the top row. The arrow that runs the length of the stack marks the scale moment. Rather than acting as a seventh principle, it is the premise on which this book opened, and crossing it is what turns each of the six from a local concern into a fleet-wide constraint.

The six principles therefore work as a decision map rather than a glossary (table 17.1). Each row links a principle to the question, metric, and practice an engineer uses when diagnosing a fleet-scale design, and to the part and chapters that develop it. The part order is an order of dependency, not a ranking of bottlenecks. The workload and its measurements determine which constraint binds, while the layered tiers in figure 17.1 locate the governing architectural band.

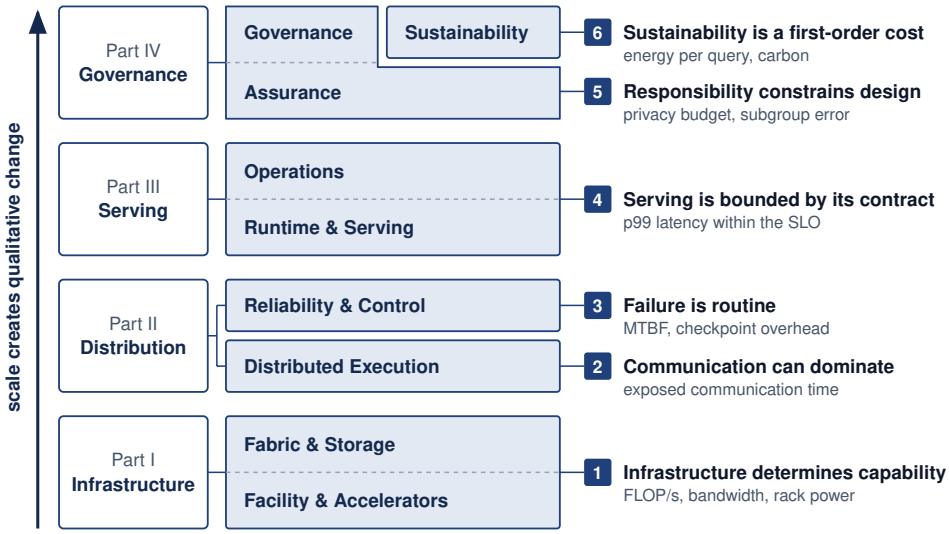

Figure 17.1: The Fleet Stack Revisited: The eight rows and four bands of figure 1.4, redrawn as the territory each of the six principles governs, with its key metric from table 17.1. The Infrastructure and Serving principles each own a whole band, and the Distribution band carries one principle per row. The Governance band divides by obligation rather than by row: Responsibility covers the Assurance row and the governance half of the top row, while Sustainability covers the resource half. The arrow runs the length of the stack because scale changes the behavior of every band.

Table 17.1: Six Principles of Distributed ML Systems: The qualitative shifts that occur when ML systems move from single machines to distributed production, grouped by the part of the book that develops them. Each principle connects to a governing question, the metric that tests it, the engineering practice it makes necessary, and the chapters where the concept is developed in depth.

Principle	Core question	Key metric	Practice it demands	Developed in
Infrastructure determines capability	What is possible?	FLOP/s, memory and bisection bandwidth, rack power	Co-design hardware with algorithms	Part I: chapter 2, chapter 3, chapter 4
Communication can dominate	What is the bottleneck?	Exposed communication time, scaling efficiency	Expect bottlenecks to migrate	Part II: chapter 5, chapter 6
Failure is routine	How do we recover?	MTBF, checkpoint overhead, rollback time	Treat failure as steady state	Part II: chapter 7, chapter 8
Serving is bounded by its contract	What reaches the user in time?	ML productivity goodput, p99 latency within the SLO, cost per request	Reserve headroom for recovery and tail latency	Part III: chapter 9, chapter 10, chapter 11, chapter 12
Responsibility constrains design	Who is affected?	Privacy budget, subgroup error rates, audit trails	Instrument first	Part IV: chapter 13, chapter 14, chapter 16
Sustainability is a first-order cost	What is the cost?	kWh per training run, energy per query, carbon	Compound small efficiency gains	Part IV: chapter 15

The first principle is that infrastructure determines capability. Physical infrastructure constrains both feasible configurations and attainable rates. Memory capacity and safe power and cooling envelopes limit how much state and active hardware a configuration can support. Cluster-wide bisection bandwidth instead limits how quickly distributed work can exchange data (principle 6); insufficient bandwidth can make a deadline or cost target unattainable without making the

computation impossible. Storage sets a further ceiling, because accelerators that the data pipeline cannot feed sit idle regardless of their peak rate (principle 7). These are the memory, network, and energy walls of the scale moment, measured at the scale of the node, the fabric, and the facility. Partitioning, batch size, and communication scheduling determine how those limits bind, and the limits themselves run from chip-level thermal density to cluster-wide bisection bandwidth.

The second principle is that communication can become the binding term at scale (Z. Jiang et al. 2024; Narayanan et al. 2021), with its cost set by the bandwidth-latency trade-off (principle 5) inside the Fleet Law's $T_{\text{comm}}(N)$ term. Communication intensity, the ratio of bytes crossing the network to local FLOPs, indicates network demand relative to arithmetic. Whether that demand dominates also depends on achieved compute and network rates, message sizes, topology, and overlap. Gradient synchronization and straggler effects must therefore be measured on the synchronous step's critical path rather than inferred from byte volume alone.

↩ The communication-intensity ratio is established in section 1.2.4.

Collective algorithms such as Ring AllReduce,[1] popularized for deep learning by Horovod (Sergeev and Balso 2018), together with gradient compression and overlapping computation with communication, address communication bottlenecks. For dense collective workloads, conventional oversubscribed fabrics often bottleneck, motivating high-bisection fabrics and transport choices matched to ML communication patterns. Recognizing when communication is the active constraint clarifies when algorithmic optimizations will help and when they merely shift work between equally constrained resources. Such a shift is the displacement of overhead (principle 14), in which reducing one term of the step increases another, and it is what makes bottlenecks migrate. Expanding the fabric across thousands of nodes also multiplies the components that a job needs at the same instant, and that multiplication makes failure routine.

[1] **Ring AllReduce**: A collective that sends $2(N-1)M/N$ bytes per worker and receives the same amount for message size M, giving $T_{\text{ring}} \approx 2(N-1)\alpha + \frac{2(N-1)}{N}\frac{M}{\beta}$ under the α-β communication model, with β denoting bandwidth in each full-duplex direction. The coefficient approaches $2\times$ message volume per worker as N grows, while remaining bandwidth-optimal for large messages.

↩ Ring collectives, communication overlap, and gradient compression are established in chapter 5.

The third principle is that failure is routine. A job that needs many components at once can encounter interruptions frequently enough that recovery becomes normal operation. The reliability gap in section 1.4.1 follows from equation 1.2. Under it, a job that needs every one of its components at once, each independently available with probability $A_{\text{component}}$, sees its availability decay exponentially in the number of components it requires. Meta's experience training Llama 3 on 16,384 GPUs, 16 times the Archetype A (GPT-4/Llama-3) reference cluster used in this chapter's worked expansion, documented 419 unexpected interruptions over 54 days, averaging one unexpected interruption every 3.1 hours (Dubey et al. 2024). A failed rank can stall every other rank in its collective. Under full-job checkpoint recovery, the affected job rolls back to its last checkpoint; other jobs in the fleet need not roll back. For a fixed amount of lost wall-clock progress, wasted accelerator-time grows with the size of the affected job, which is why checkpoint cadence and recovery architecture are first-order design dimensions.

Architects must embed failure handling from the beginning. Checkpointing strategies balance write overhead against expected lost work, and the Young-Daly interval (principle 13) estimates that balance under its failure and checkpoint-cost assumptions. Elastic training supports changing cluster membership, while graceful degradation preserves a defined level of service as capacity diminishes. Each mechanism needs a tested recovery contract; its name alone does not establish that progress or service quality survives a failure. Recovery should automate known failure paths and escalate cases whose safe response requires investigation. Recovery also needs capacity. A restarted synchronous job makes progress only when the scheduler can again grant it a gang allocation, all of its accelerators at once and in a feasible topology (principle 15). That is why fleet orchestration keeps spare capacity, preempts lower-priority work, and admits only allocations a job can actually use. Coordination, the third C^3 axis, is not confined to this principle. It is paid in recovery and gang scheduling here, in request routing for serving, and in the approval gates of governance.

↩ Checkpoint cadence, mean time between failures (MTBF) analysis, and elastic recovery protocols are established in chapter 7.

Communication and failure together constitute the operational reality of the Distribution band. Deployment carries both into a setting where the unit of work is a user request rather than a training step, which gives the fourth principle, that serving is bounded by its contract. The accounting begins before any request arrives. The ML goodput decomposition (principle 16) assigns each interval of capacity that the built fleet loses to exactly one of three factors, so that every loss has one owner. Scheduling goodput is the share of capacity held by complete allocations, runtime goodput is the share of that allocated time retained as progress, and program goodput is the fraction of the ideal rate that the retained steps achieve. A serving fleet's useful capacity is then the work it completes

within its latency, availability, and cost targets, not the work its accelerators execute in isolation, which is what section 10.3 measures as service-level objective (SLO) goodput.

Production inference that waits for every worker in a fan-out inherits tail effects (Dean and Barroso 2013a), because completion depends on the slowest required response. A request that touches 100 servers has a 63.4 percent chance that at least one of them responds beyond its 99th-percentile latency when server delays are independent,[2] so a latency that is rare for one server becomes the common case for the request. Routing is where coordination enters the request path. Probing two of R replicas and sending the request to the less loaded (principle 19) sharply shortens the longest queue in the idealized balls-into-bins model, at the cost of only two probes. On live traffic the gain still depends on arrivals, service times, and how fresh the load estimate is, so it narrows the tail rather than bounding every request's latency.

Inside each replica, the memory wall returns as a serving constraint. Compute (TFLOP/s) can be abundant while memory bandwidth (GB/s) gates small-batch autoregressive decode (principle 18), because weight and KV-cache traffic from high-bandwidth memory (HBM) leaves compute units underused. Batching amortizes each weight read across the requests in the batch, so the binding constraint can shift toward compute, KV traffic, or the latency budget. A replica's execution envelope must therefore be measured before any fleet-level policy is chosen. The contract also has a lifetime. A heavily used model can spend more on inference than it did on training (principle 17), so the cost per request, not the one-time training bill, eventually dominates what the model costs to own.

Beyond the data center, the contract takes other forms. On a phone-class device the contract is the device's own power, thermal, and foreground-latency envelope, and the device enforces it by admitting background work, including local adaptation that updates the model on the device's own data, only when the envelope has room. That admission policy has a consequence that reaches back into training. Only devices whose policy accepts the work take part in adaptation, so the serving-side contract also decides whose data the shared model learns from (principle 20), a bias the fifth principle must then account for. Across a platform of hundreds of models, staged rollouts and per-service reliability budgets enforce the contract release by release, so every release must carry evidence that its targets still hold. Once the fleet can deliver a service, governance determines whether it is allowed to.

The fifth principle is that responsibility constrains design. Responsible AI translates risk-management frameworks into engineering constraints (Tabassi 2023). Safety-specific risks such as reward hacking, distribution shift, and scalable oversight show why those constraints need implementation mechanisms (Amodei et al. 2016). Reward hacking occurs when a system games its measured objective, distribution shift when deployment inputs differ from the training data, and scalable oversight is the difficulty of supervising behavior that people cannot check directly. Fairness, transparency, accountability, privacy, and safety are first-class requirements that shape system architecture throughout the ML lifecycle. The part-level principles behind this one span provenance, privacy, robustness, and feedback. A model can be trusted only as far as the verified record of its data, code, and weights reaches (principle 22). Released outputs can disclose what a model learned (principle 23). Robustness evaluation and defenses require mechanism-specific resource budgets (principle 24). Deployed decisions can shape future training observations (principle 27).

The fairness impossibility theorem (principle 26) shows why some requirements cannot be satisfied by tuning alone. With unequal base rates and imperfect prediction, a risk score cannot be both calibrated within each group and balanced across groups within each outcome class (Kleinberg et al. 2017). Calibration requires that, within every group, each score equal the observed rate of the outcome among the people who receive it, and balance requires that people with the same outcome receive the same average score in every group. This is the score-level counterpart of the conflict between equal opportunity and demographic parity in the admissions example of section 16.5.1.5. The deployment therefore needs justified criteria, evidence that they hold, and an owner authorized to respond when they fail. Monitoring exposes supported disparities but cannot choose the criterion or provide a remedy by itself. These obligations belong in the design alongside performance requirements, which is why the practice this principle demands is to instrument first. Subgroup error rates, privacy accounting, and audit trails must be collected before the decision they are meant to justify, because a disparity that was never measured cannot be detected, attributed, or remedied after deployment.

[2] **Tail latency**: The estimate is one minus the probability that every server answers within its own percentile, which under independence is the per-server probability raised to the number of servers; section 10.11 analyzes the tail effects it summarizes. Shared congestion correlates server delays, so a real fleet can sit above or below this figure, and the independence assumption has to be checked against measured traces.

↩ The power of two choices and the queue-length bounds behind it are established in section 10.11.3.

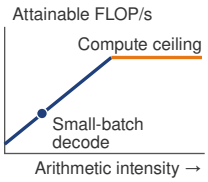

Small-batch decode can sit left of the roofline ridge.

↩ On-device adaptation within phone-class power and thermal envelopes is established in section 11.3.5.

↩ Reward hacking, scalable oversight, and distribution shift as concrete AI safety problems are established in section 16.8.3.

↩ The calibration and score-balance conditions that the fairness impossibility theorem relates are established in section 16.5.1.5.

The sixth principle is that sustainability is a first-order cost. Environmental impact of large-scale ML elevates resource efficiency to a primary engineering constraint (Strubell et al. 2019; Patterson et al. 2021). The energy-scale invariant of section 1.5.5 supplies the accounting discipline, a fixed boundary for useful work and for the energy it consumes. The Fleet Law explains why exposed communication and synchronization cost joules as well as seconds, since a step's energy is approximately its average power times $T_{\text{step}}(N)$. The Jevons paradox of AI adapts Jevons's resource-efficiency argument (Jevons 1865) to explain why hardware efficiency gains alone do not eliminate power and carbon costs (principle 25). Cheaper computation can expand deployment, so usage growth can offset per-query savings. An absolute carbon budget bounds emissions only when accounting and admission decisions enforce it over the stated lifecycle and time horizon.

↩ Carbon accounting, energy metrics, and Jevons rebound effects are established in chapter 15.

Sustainability thus becomes an engineering discipline with its own accounting chain, in which IT energy scaled by the facility's power usage effectiveness and the grid's carbon intensity gives operational emissions, and the embodied carbon of manufacturing the hardware joins the total. Two links show how far the chain can swing. Training GPT-3 consumed about 1,287 MWh, roughly 120 US household-years of electricity (Li 2020), and the same energy drawn from Poland's coal-heavy grid emits about 41× the carbon it would on Quebec's hydroelectric grid. The chain also explains the practice this principle demands. Small efficiency gains compound across every accelerator and every query in the fleet, but usage growth can absorb them, so compounding efficiency bounds total emissions only when an absolute budget caps the total. Carbon-aware scheduling, lifecycle assessment,[3] and efficiency optimization are the engineering means of meeting that budget.

[3] **Lifecycle assessment**: Section 15.8 evaluates environmental impact across a system's lifespan, including embodied carbon from hardware manufacturing. Embodied emissions can become a substantial share of total impact on clean grids or when hardware is lightly used.

↩ The operational and embodied carbon calculation is established in section 15.5.

Taken together, the six principles describe the qualitative change that scale produces across a fleet. Each larger regime brings behavior that emerges from familiar mechanisms, as the scale moment first showed. A training job that scales from 8 GPUs to 8,000 can meet communication bottlenecks, load imbalance, or synchronization overhead that were negligible at 8, and a serving fleet handling a million requests a day expects a one-in-a-million per-request failure about once a day. Stragglers that bottleneck clusters, network partitions that split training, and heterogeneity across hardware generations that complicates load balancing therefore become routine design inputs rather than anomalies. Because empirical scaling laws (principle 9) keep carrying teams that pursue lower loss into larger regimes, a production design must weigh the quality gain those laws predict against the full-stack costs the six principles govern.

17.3 The Complete Production System

A team that doubles its training cluster to chase the next point on a scaling curve halves the ideal arithmetic time per step, yet the added accelerators also increase collective traffic, shorten the interval between hardware interruptions, and draw against a fixed facility power budget. Empirical scaling laws set the demand for that fleet growth and the Fleet Law (equation 17.1) sets its price, so together they say when growth pays and where it breaks. The scaling curve estimates what growth buys; it does not oblige an organization to buy it. Scaling pays only while the predicted quality gain justifies what the fleet charges across every term that moves with N. Increasing N reduces the ideal compute term T_{compute}/N for fixed work and unchanged per-device rates, whereas $T_{\text{comm}}(N)$, $T_{\text{sync}}(N)$, and interruption exposure depend on the partition, schedule, and component count and carry no such guarantee. Within a fixed facility, the power and cooling envelope does not grow with the job, and a successful model then pays its serving cost for as long as it runs.

Growth can therefore break on either side. On the model side, the breakdown conditions of table 1.2 mark where extrapolating a fitted quality curve becomes unreliable, for example when model capacity outgrows the available tokens or added data stops supplying new information. On the systems side, the scaling wall of section 6.4.1 marks where added communication and coordination consume the compute time that added accelerators save (principle 12), while deliverable site power caps the configuration regardless of how work is partitioned. A complete production design has to evaluate both sides together.

Consider Archetype A (GPT-4/Llama-3). Because GPT-4's configuration is not published, its worked examples in this book use a GPT-3-class dense transformer, trained with mixed precision on a reference cluster of 1,024 GPUs and then shipped into a global serving fleet. During training, checkpoint bursts must land without starving the data loaders, while collective results must arrive before the computations that depend on them. Communication can overlap independent work, but

its exposed remainder delays progress. At that scale, recovery planning uses observed write costs and failure rates, with the Young-Daly interval supplying an approximation where its assumptions hold (Young 1974; Daly 2006). Serving SLOs require headroom for rollback and fallback, while governance controls enforce approved conditions on which data, models, and outputs may reach users. These are coupled rates, budgets, failure exposures, and policy constraints that must be designed before deployment.

The production system begins with physical capacity, but capacity matters only when the rates line up. Storage, fabric, and compute must therefore be co-designed so that data ingestion meets accelerator consumption, checkpoint writes fit the recovery budget, and the fabric carries the actual collective traffic. Matching these distinct rates to workload demand matters more than making their raw bandwidths equal.

Parallelization strategies must fit those physical rates and the expected failure pattern. Data parallelism, tensor parallelism, and pipeline parallelism are not interchangeable templates; each changes memory pressure, communication volume, checkpointing behavior, and scheduler complexity. Hybrid strategies become useful only when the surrounding storage, fabric, and recovery design can support the shape of the partition.

↩ How data, tensor, pipeline, and hybrid partitioning each move memory pressure, communication volume, and the pipeline bubble is established in section 6.8.

Serving adds a separate acceptance boundary. The contract of the fourth principle judges the model by its latency, availability, and cost on live traffic, whether it runs in a data center replica or on a phone-class device. Some security and robustness controls execute on the request path; others operate during training, release review, or asynchronous monitoring. Each needs an evidence and response path suited to its obligation. Training throughput therefore does not establish that a model will meet live latency, recovery, security, subgroup, and carbon requirements. A fleet expansion raises two linked decisions, whether to accept the larger training job and whether to release the model it produces, and each needs its own evidence.

To evaluate the first of those decisions on the Archetype A baseline, suppose a proposed expansion doubles the reference cluster to 2,048 GPUs while retaining the fabric and checkpoint tier. Because doubling N halves the ideal T_{compute}/N term, the first measurement is whether that saving survives the resulting changes in $T_{\text{comm}}(N)$ and $T_{\text{sync}}(N)$ and their overlap. The recovery side can be estimated before any measurement. Assume independent GPU failures at an identical constant rate, with each failure interrupting the job. Counting those failures alone at an assumed mean time to failure (MTTF) of 50,000 hours per GPU, the job's mean time between interruptions is that MTTF divided by N, so it falls from 48.8 hours to 24.4 hours.

At the Young-Daly interval of equation 7.7, the checkpoint-and-rework overhead is about $\sqrt{2T_{\text{write}}/\text{MTBF}_{\text{system}}}$ of wall-clock time. Because the job's mean time between interruptions falls in inverse proportion to N, that overhead grows as $\sqrt{N}$ (section 7.8.1), so doubling the cluster multiplies it by about 1.41×. The model's 2.45 TB resumable checkpoint, which stores its weights and optimizer state, takes about 24.5 seconds to write at the assumed checkpoint write bandwidth, which puts the overhead at about 1.7 percent of wall-clock time on the reference cluster and about 2.4 percent on the expanded one.

These figures are optimistic. They count GPU failures alone, while switches, network interface cards, and software add interruptions of their own. The full-system analysis in section 7.8.2, which adds PCIe-switch and network-interface failures for a cluster of 10,000 GPUs, arrives at a correspondingly larger overhead from the same formula. Write contention with data loading or a changed checkpoint size would also require revising them. Appendix G lists the parameter count, checkpoint bytes per parameter, GPU MTTF, and checkpoint write bandwidth behind these figures, so a reader can recompute them or substitute other values.

These estimates still yield a decision rule. In this example, measured step time covers uninterrupted execution and excludes the checkpoint-and-rework overhead, so that overhead, which runtime goodput counts as lost time (section 9.5), enters the run-time estimate as a separate factor. A step time measured over the whole run already includes recovery pauses and must not take the factor again. Because only the fraction $1-\text{overhead}$ of each wall-clock hour produces retained training progress, total run time for fixed work is proportional to $T_{\text{step}}(N)/(1-\text{overhead})$. Doubling the cluster raises the checkpoint-and-rework overhead from 1.7 percent to 2.4 percent, which leaves a slightly smaller productive fraction of each hour on the expanded cluster. To finish in less wall-clock time, the

expanded cluster's uninterrupted step-time speedup must therefore exceed the ratio of the reference productive fraction to the expanded productive fraction, $(1-\text{overhead}_{\text{ref}})/(1-\text{overhead}_{\text{exp}})$. Here that break-even step speedup is about 1.01×, so recovery overhead alone does not block the expansion.

Short of a superlinear step speedup, however, the expansion cannot save accelerator-hours. Total accelerator-hours equal the GPU count times wall-clock run time. Even if uninterrupted step time halves at the ideal $2\times$ speedup, the higher checkpoint-and-rework overhead keeps the expanded cluster running for slightly more than half the original wall-clock duration, so doubling the GPUs multiplies total accelerator-hours by the same ratio of productive fractions that set the break-even speedup, an increase of about 0.7 percent. Every point of scaling efficiency lost to exposed communication raises that bill further. The expansion therefore buys calendar time with accelerator-hours and energy, and the communication measurement, not the recovery estimate, sets that price.

That price buys time, not quality, because the work is fixed. Quality enters the comparison only if the doubled cluster spends twice the compute instead of finishing sooner, and there the fitted scaling curve prices the gain. In the compute power law that Kaplan et al. (2020) fitted (charted in the scaling laws of section 1.3, figure 1.8), loss falls as compute raised to the power -0.050, so doubling compute multiplies loss by $2^{-0.050} \approx 0.966$, a reduction of about 3.4 percent. A frontier run lies beyond the range that fit studied, so the figure shows the shape of the trade rather than this run's gain. Each doubling buys the same fractional reduction, and therefore a smaller absolute one than the doubling before it, while its price, twice the accelerator-hours plus the recovery and communication overhead above, is paid in full every time.

The release decision then applies its own test to the model the expanded run produces. An end-to-end evaluation contract (principle 21) fixes that test before exposure. The contract is one versioned release record that binds the candidate to production-equivalent measurements, accountable acceptance authority, tested recovery, and explicit conditions that invalidate the evidence. Its serving gate requires the model to meet its p99 latency target under production-equivalent load with headroom left for rollback and fallback. For Archetype A, the governance gates require a provenance record that pins the data, code, and weights, evidence bounding what released outputs can disclose, and training-energy accounting over a stated boundary. A failing gate sends the revision back to the fleet-stack row that owns the failure, which may mean changing the partition, placement, capacity allocation, or release plan, and the release proceeds only after the revised candidate passes.

The constraints behind those revisions can conflict. A communication optimization may increase failure exposure, sustainability requirements may restrict infrastructure choices, and governance controls may consume serving latency or capacity. The displacement of overhead requires measuring those transferred costs before accepting a local speedup.

↩ The end-to-end evaluation contract and its versioned release record are established in section 12.5.8.

Archetype A traces only one path through these cross-layer trade-offs; taken together, the three lighthouse archetypes show that no single corner of the C^3 taxonomy dominates. Archetype B (DLRM at Scale) is a recommendation model whose embedding tables are sharded across many nodes, and Archetype C (Federated MobileNet) is a small vision model trained across millions of phones that keep their data local. Each workload was first characterized by a representative binding term, and following each one up the fleet stack shows what each band, and the part of the book that develops it, added to that first characterization (table 17.2).

↩ The three lighthouse archetypes, and whether each stresses compute, communication, or coordination, are established in section 1.6.1.

Table 17.2: Archetypes across the Fleet Stack: The book's three running workloads traced through the four bands of the fleet stack, one band per part, beneath the representative binding C^3 term with which the introduction characterized each. The same Fleet Law applies to all three; the regimes shown here expose different constraints, not permanent bottleneck assignments.

Fleet stack band	Archetype A (GPT-4/Llama-3)	Archetype B (DLRM at Scale)	Archetype C (Federated MobileNet)
Representative C^3 term	Communication, fleet-wide	Coordination, in routing and placement	Local compute, within the device envelope
Infrastructure (Part I)	Training state and checkpoints far exceed one accelerator's memory, motivating partitioning; frequent tensor-parallel collectives favor the fastest interconnect domain	Multi-terabyte embedding tables exceed any device's memory, and live feature lookups demand random-access storage	Phone-class power, memory, and thermal envelopes cap what local training can attempt

Fleet stack band	Archetype A (GPT-4/Llama-3)	Archetype B (DLRM at Scale)	Archetype C (Federated MobileNet)
Distribution (Part II)	3D parallelism and AllReduce must fit the step budget, and Young-Daly checkpoint cadence tracks interruptions measured in hours	Embedding sharding turns lookups into AllToAll exchanges whose hottest receiver sets completion time	Federated averaging replaces AllReduce, and over-selection with round deadlines bounds the wait for intermittent clients
Serving (Part III)	Small-batch decode is bound by memory bandwidth, so batching, KV-cache policy, and lifetime serving cost shape deployment	Feature-parallel sharding and routing must hold tail latency within budget at millions of queries per second	Adaptation runs only in admitted windows that preserve the foreground deadline, and rollouts span weeks of version skew
Governance (Part IV)	Data provenance, output leakage, and the energy of each training run gate release	Recommendations shape future training data, requiring fairness and drift evidence with accountable response across the serving lifecycle	Secure aggregation limits visibility of individual updates; differential privacy bounds released information under stated accounting assumptions; participation bias shapes whose data contribute

Each column traces the constraints one workload encounters across the bands, while each row compares how a single band constrains different workloads. These constraints need not all bind at once. The representative terms describe operating regimes, not permanent assignments of archetypes to C^3 terms. Federated MobileNet can be device-compute-bound during local training yet communication- or coordination-bound during a global round with intermittent clients. The same six principles apply to all three; the engineer must measure the workload at the relevant boundary before choosing which term to reduce. Because the binding term must be found rather than assumed, the book's method reduces to a diagnostic procedure, a sequence of steps that is distinct from the principles it applies:

1. Start with the observed symptom, the user-visible or operator-visible failure.
2. Attach the metric that makes the symptom falsifiable.
3. Map the symptom to C^3 by identifying whether compute, communication, or coordination binds first.
4. Locate the fleet-stack row that owns the intervention.
5. State the cost that the proposed fix displaces elsewhere.
6. Preserve the audit, safety, security, or responsibility evidence needed to justify the decision.

On the expanded cluster, a training run that falls behind its schedule exercises every step of that procedure. The schedule lag is the symptom, and the falsifiable metric is model FLOPs utilization measured per step, the measure of program goodput in section 9.2.2, which has dropped since the expansion. A step profile shows accelerators idle at gradient-synchronization barriers, but that wait alone does not distinguish slow collective transfer from ranks arriving late after computation, data loading, or scheduling delays. Aligned per-rank traces must therefore separate arrival skew from exposed collective completion time. When those traces show little arrival skew and increased exposed transfer time, they point to communication as the binding term and to the distributed-execution row, where the collective and its overlap are scheduled, as the owner of the intervention. A proposed fix such as gradient compression shrinks the exchanged bytes while displacing cost into convergence and accuracy risk, which means the evidence to preserve is the convergence and evaluation record showing that the compressed run still reaches its quality target.

The same procedure redirects other symptoms. A fairness alert may point to missing labels, insufficient review capacity, or a serving-path threshold policy rather than model weights alone. A security incident may point to artifact provenance, access permissions, or rollback evidence rather than network perimeter failure. The discipline is to follow the constraint until the row that owns it and the trade-off it carries are visible. When the fix changes what ships, the evidence from the final step enters the release record of the same end-to-end evaluation contract.

17.4 Competencies Mastered

The proposed cluster expansion tests competence across the whole stack. Accepting it requires evidence that reduced compute time survives communication and recovery costs, while releasing its model requires a separate serving and governance judgment. Mastery means identifying those dependencies, choosing the measurements that test them, and defending the decision within the stated assumptions. Each band of the stack contributes one checkable part of that defense.

In the Infrastructure band, the governing question is which infrastructure limit binds the workload. An engineer who has mastered this material can match accelerator memory and compute to the model, provision storage and network bandwidth for the required rates, and account for power and cooling as capacity constraints rather than facilities details. In the expansion, that skill decides whether the retained fabric, checkpoint tier, and site power can carry 2,048 GPUs at all.

In the Distribution band, the governing question is how to make a workload larger than any single machine behave as one coordinated system. An engineer can analyze communication patterns, select a parallelization strategy, and estimate failure cadence from component MTTF, fleet size, and the failure-dependence model. Given a collective's cost model, such as the Ring AllReduce cost recalled in section 17.2, the engineer can also predict which of its terms grows with N. For the ring, the latency term grows linearly with the worker count, while the bandwidth term levels off as each worker's traffic approaches twice the message size. In the expansion, this band produced the recovery estimate and the break-even speedup, and it identified exposed communication as the measurement that sets the price.

In the Serving band, the governing question is how a trained model keeps serving value under live traffic. An engineer can budget the serving tax, select batching and placement policies, and monitor system performance and input distributions in production so that a working model does not become an inefficient or stale service. In the expansion, this band supplies the latency gate that the released model must pass with headroom left for recovery.

In the Governance band, the governing question is who bears the costs and risks of the system's decisions. An engineer can translate security, privacy, fairness, robustness, and sustainability obligations into controls and evidence: access boundaries, privacy budgets, subgroup evaluations, adversarial, out-of-distribution, and concept-drift tests, incident procedures, and resource accounting. In the expansion, this band supplies the security, subgroup, and carbon evidence that the release decision requires.

↩ The serving tax that distributed execution adds to each request's latency budget is established in section 10.3.

These competencies connect model development to the consequences of deployment. They also supply a test of transfer. When a new workload or substrate appears, the engineer must identify its accounting boundary, constraints, and evidence requirements before accepting an efficiency claim.

17.5 Where the Binding Constraint Moves Next

When a reasoning model spends seconds of test-time compute on a single user prompt or an optical interconnect replaces copper across a switching tier, that test of transfer applies directly. The binding constraint shifts again. The same six principles govern workloads and substrates that are only now emerging. Test-time compute, introduced among the scaling regimes in section 1.3.3.1, already moves part of the scaling curve from the training cluster into the serving path, where section 10.6.9 prices it as a per-request compute budget. Beneath the fleet, new interconnect and packaging substrates change the rates that every Fleet Law term inherits. Each shift relocates cost without exempting it from the Fleet Law.

Gains from scheduling, placement, and overlap depend on a physical fabric, so software and substrate choices must be evaluated together. Once software has removed avoidable coordination, the remaining frontier is the substrate that moves tokens, activations, and power through the fleet. For Archetype A, the pressure on that substrate comes from specific collective patterns rather than from scale in general. Tensor parallelism issues frequent collectives across partitioned tensors, and in mixture-of-experts (MoE) variants of such models, routing issues an AllToAll dispatch at each expert layer. Across thousands of accelerators, link energy accumulates with the volume of that collective traffic, which makes lower-energy interconnects a plausible fabric-efficiency lever.

Candidate substrates attack different factors of that physical account: energy per bit, path length, bandwidth density, or the distance between memory and compute. Optical I/O and co-packaged optics lower the energy required to move each bit across the fabric, while 3D integration extends the

vertical stacking of HBM and advanced packaging to shorten interconnect paths. Neuromorphic event-driven substrates instead organize computation around sparse events such as spikes and keep state beside the compute units that use it. Automated architecture search and AI-for-systems optimization offer ways to explore candidate designs, but each proposal still needs a workload-specific comparison of useful work, data movement, power, and implementation cost.

↩ Neuromorphic and other non-von Neumann directions are established in section 15.10.5.

A lower energy per bit reaches the fleet only in proportion to the links' share of fleet power, the affected-share arithmetic of Amdahl's law applied to energy. At fixed traffic, halving the energy per bit saves at most half of the links' share of fleet power, not half of fleet power. A lower link-energy measurement alone therefore does not establish a better fleet. The optical dividend of section 3.10.2 sizes the optics power saving for one switch and for a cluster of switches, the starting point for that share calculation. Even a substrate that wins the energy comparison still needs fault tolerance and responsible governance before it improves the fleet.

Whatever the substrate, the discipline stays the same. An engineer chooses the unit of useful work, accounts for its full resource cost, and states which comparisons the evidence supports. Narrowing that boundary to a single component, layer, or benchmark in isolation produces the recurring failure modes of fleet design.

17.6 Fallacies and Pitfalls

The closing mistakes all come from optimizing the wrong unit of design. A model can improve while the fleet becomes slower, less reliable, more expensive, or harder to govern.

Fallacy: *A faster model is automatically a better system.*

Raw model speed matters, but it is only one term in the Fleet Law. In the illustrative equal-work, barrier-dominated step, a worker running at 80 percent of peer speed reduces throughput by 20 percent relative to uniformly fast workers. A faster model, delivered as a faster kernel on every worker but that one, leaves the step bound by the straggler, because speeding the workers already waiting at the barrier does not shorten the limiting computation. A design that reduces kernel time while increasing checkpoint pressure, network contention, serving variance, or governance risk can make the total system worse. The systems question is whether the change improves useful throughput, the ML goodput of section 9.2, under the actual constraints: data movement, synchronization, failure recovery, power, latency, and accountability.

Pitfall: *Optimizing one layer while hiding the constraint it creates in another.*

A local improvement becomes dangerous when it moves cost to a layer that no one is measuring. Larger batches may raise accelerator utilization while worsening tail latency, aggressive compression may reduce communication while increasing accuracy risk, and carbon-aware scheduling may lower emissions while requiring more slack in the service budget. Good engineering respects the displacement of overhead, keeping the transferred constraint visible so the system can choose deliberately rather than inheriting a hidden bottleneck.

Fallacy: *Fleet-scale lessons are tied to today's hardware and software stack.*

Specific processors, frameworks, and serving engines will change, but the durable relationships remain. Data must reach compute at the required rate. Workers must communicate before synchronization stalls the job. Failures must be expected at large component counts. Serving systems must meet tail-latency budgets. Governance constraints must be engineered into the operating path. The technology names are examples; the rate, failure, and accountability relationships are the lesson.

Pitfall: *Treating governance and sustainability as external reviews rather than system constraints.*

Postponing security, fairness, privacy, and carbon accounting can create architectural debt. Retrofitting audit trails, demographic monitoring, access controls, deletion workflows, or carbon-aware placement may require reconstructing evidence or redesigning deployed components. Identifying these obligations during design lets teams budget the required paths and test whether they work before release; the actual cost still depends on the control and deployment.

17.7 Summary

At fleet scale, improving a model or accelerator does not establish that the complete service improves. The engineering object is the coupled system that moves data, schedules work, survives failure, serves users, and satisfies governance obligations under physical and organizational constraints. Within their fitted regime, scaling laws set the demand for fleet growth, the Fleet Law sets its price, and the six principles locate where in the fleet that price is paid. The method that recurs across the book follows from all three. It identifies the binding constraint, quantifies its cost, and evaluates the proposed change across every affected operating path, so that the cost it displaces stays visible instead of resurfacing unmeasured in another layer.

Key Takeaways: The fleet is the unit of design

- **The fleet is the object**: The six principles reduce to one habit, following the binding constraint across infrastructure, communication, coordination, serving, and governance. The fleet, not any single model or component, is the unit to build, measure, and optimize.
- **Growth must earn its price**: Within a fitted regime, each doubling of compute buys a smaller absolute loss reduction, while interruption exposure and checkpoint overhead grow with the fleet and exposed communication can grow with it. Expansion is a decision justified by measurements and release evidence, not a default.
- **Scale changes the probability model**: When server delays are independent, a request that touches 100 servers has a 63.4 percent chance that at least one responds beyond its 99th-percentile latency, and a job that needs every one of N independent components sees its interval between interruptions shrink as $1/N$. Fleet behavior is not single-node behavior repeated.
- **Orchestration gains require system measurements**: Scheduling, placement, and communication overlap improve useful throughput only when their gains survive recovery costs, serving deadlines, and resource budgets. Measure the complete operating path before accepting a local improvement as a fleet-wide gain.
- **Obligations belong in the path**: Security, privacy, fairness, robustness, carbon, and auditability are production constraints, not external reviews. The discipline is to design them into data pipelines, schedulers, serving paths, and operating procedures before scale makes the trade-off irreversible.

Closing Distillation: The discipline of ML systems

Section 1.1 traced the scale moment through a training run that stopped every few hours. Meta's Llama 3 cluster of 16,384 GPUs recorded 419 unexpected interruptions in 54 days and still kept more than 90 percent of elapsed time as useful training (Dubey et al. 2024). Those numbers now read as engineering. An unexpected interruption every 3.1 hours is the reliability gap of a job that needs every one of its accelerators at once, and it measures the run's recovery exposure. The effective training fraction, the run's runtime goodput, measures how much progress the recovery machinery retained across elapsed time. The report credits faster startup, checkpointing, and diagnosis, including communication traces that made stalled collectives inspectable, for that retained fraction, so the fraction summarizes all of them together rather than measuring any single Fleet Law term. Each worker still computed, exchanged, and waited, as in the synchronous step with which the book began; Compute, Communication, and Coordination are those three activities, measured.

A proposed larger run can now be judged rather than presumed necessary. Under the expansion example's independent, identical constant-rate failure model and unchanged checkpoint write time, doubling the fleet halves the expected interruption interval and raises optimized checkpoint-and-rework overhead by about 1.41×. It buys calendar time only by spending more

> accelerator-hours on the same work. Whether that run should be built depends on the shrinking quality gain its fitted scaling curve predicts for the doubling, the measured time and energy required to deliver it, and the evidence its release requires. The engineering achievement is a justified decision, including the decision to retain the current fleet when expansion cannot meet those conditions. The same discipline applies whether the work scales training compute across thousands of accelerators or serves a model to millions of users under a latency contract. The engineer follows the binding constraint across the fleet, accounts for every displaced cost, and accepts a design only when the evidence supports it.

Prof. Vijay Janapa Reddi, Harvard University

REFERENCES

17

References

Abadi, M., A. Chu, I. Goodfellow, H. B. McMahan, I. Mironov, K. Talwar, and L. Zhang. 2016. "Deep Learning with Differential Privacy." *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, CCS '16, 308–18. https://doi.org/10.1145/2976749.2978318.

Adi, Yossi, Carsten Baum, Moustapha Cisse, Benny Pinkas, and Joseph Keshet. 2018. "Turning Your Weakness into a Strength: Watermarking Deep Neural Networks by Backdooring." *27th USENIX Security Symposium (USENIX Security 18)*, 1615–31.

Advanced Micro Devices. 2021. *Introducing AMD CDNA 2 Architecture*.

Agrawal, Amey, Nitin Kedia, Ashish Panwar, Jayashree Mohan, Nipun Kwatra, Bhargav S. Gulavani, Alexey Tumanov, and Ramachandran Ramjee. 2024. "Taming Throughput-Latency Tradeoff in LLM Inference with Sarathi-Serve." *18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24)*, 117–34.

AI, Meta. 2022. *Building the Most Powerful AI Supercomputer: Meta's AI Research SuperCluster*. Meta AI Blog.

Ainslie, Joshua, James Lee-Thorp, Michiel de Jong, Yury Zemlyanskiy, Federico Lebron, and Sumit Sanghai. 2023. "GQA: Training Generalized Multi-Query Transformer Models from Multi-Head Checkpoints." *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 4895–901. https://doi.org/10.18653/v1/2023.emnlp-main.298.

Aizman, Alex, Gavin Maltby, and Thomas Breuel. 2019. "High Performance I/O for Large Scale Deep Learning." *2019 IEEE International Conference on Big Data (Big Data)*, 5965–67. https://doi.org/10.1109/bigdata47090.2019.9005703.

Aji, Alham Fikri, and Kenneth Heafield. 2017. "Sparse Communication for Distributed Gradient Descent." *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, 440–45. https://doi.org/10.18653/v1/d17-1045.

Akidau, Tyler, Robert Bradshaw, Craig Chambers, Slava Chernyak, Rafael J. Fernández-Moctezuma, Reuven Lax, Sam McVeety, et al. 2015. "The Dataflow Model: A Practical Approach to Balancing Correctness, Latency, and Cost in Massive-Scale, Unbounded, Out-of-Order Data Processing." *Proceedings of the VLDB Endowment* 8 (12): 1792–803. https://doi.org/10.14778/2824032.2824076.

Alexandrov, Albert D., Mihai F. Ionescu, Klaus E. Schauser, and Chris J. Scheiman. 1995. "LogGP: Incorporating Long Messages into the LogP Model—One Step Closer Towards a Realistic Model for Parallel Computation." *Proceedings of the Seventh Annual ACM Symposium on Parallel Algorithms and Architectures - SPAA '95* 30: 95–105. https://doi.org/10.1145/215399.215427.

Al-Fares, Mohammad, Alexander Loukissas, and Amin Vahdat. 2008. "A Scalable, Commodity Data Center Network Architecture." *ACM SIGCOMM Computer Communication Review* 38 (4): 63–74. https://doi.org/10.1145/1402946.1402967.

Alistarh, Dan, Demjan Grubic, Jerry Z. Li, Ryota Tomioka, and Milan Vojnovic. 2017. "QSGD: Communication-Efficient SGD via Gradient Quantization and Encoding." *Advances in Neural Information Processing Systems (NeurIPS)* 30: 1709–20.

Alizadeh, Mohammad, Albert Greenberg, David A. Maltz, Jitendra Padhye, Parveen Patel, Balaji Prabhakar, Sudipta Sengupta, and Murari Sridharan. 2010. "Data Center TCP (DCTCP)." *Proceedings of the ACM SIGCOMM 2010 Conference*, 63–74. https://doi.org/10.1145/1851182.1851192.

Amazon Web Services. 2017. *Summary of the Amazon S3 Service Disruption in the Northern Virginia (US-EAST-1) Region*. AWS service event summary.

Amazon Web Services. 2020. *Amazon S3 Update – Strong Read-After-Write Consistency*. AWS News Blog.

Amazon Web Services. 2026a. *Checking Object Integrity in Amazon S3*. Amazon S3 User Guide.

Amazon Web Services. 2026b. *Managing the Lifecycle of Objects*. Amazon S3 User Guide.

Amazon Web Services. 2026c. *Naming Amazon S3 Objects*. Amazon S3 User Guide.

Amazon Web Services. 2026d. *Understanding Archive Retrieval Options*. Amazon S3 User Guide.

AMD. 2023. *AMD Instinct MI300X Accelerators*. AMD product documentation.

Amdahl, Gene M. 1967. "Validity of the Single Processor Approach to Achieving Large Scale Computing Capabilities." *Proceedings of the April 18-20, 1967, Spring Joint Computer Conference on - AFIPS '67 (Spring)*, AFIPS '67 (spring), 483–85. https://doi.org/10.1145/1465482.1465560.

Amershi, Saleema, Andrew Begel, Christian Bird, Robert DeLine, Harald Gall, Ece Kamar, Nachiappan Nagappan, Besmira Nushi, and Thomas Zimmermann. 2019. "Software Engineering for Machine Learning: A Case Study." *2019 IEEE/ACM 41st International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP)*, 291–300. https://doi.org/10.1109/icse-seip.2019.00042.

Amiel, Frederic, Christophe Clavier, and Michael Tunstall. 2006. "Fault Analysis of DPA-Resistant Algorithms." In *Fault Diagnosis and Tolerance in Cryptography*. Springer Berlin Heidelberg. https://doi.org/10.1007/11889700_20.

Amodei, Dario, and Danny Hernandez. 2018. "AI and Compute." *OpenAI Blog* 2.

Amodei, D., C. Olah, J. Steinhardt, P. Christiano, J. Schulman, and D. Mané. 2016. "Concrete Problems in AI Safety." *arXiv Preprint arXiv:1606.06565*, ahead of print. https://doi.org/10.48550/arXiv.1606.06565.

Angwin, Julia, Jeff Larson, Surya Mattu, and Lauren Kirchner. 2016. *Machine Bias*. ProPublica investigation.

Anslow, Pete. 2016. *RS(544,514) FEC Performance Including Precoding*. IEEE P802.3cd Task Force.

Anthony, L. F. W., B. Kanding, and R. Selvan. 2020. *Carbontracker: Tracking and Predicting the Carbon Footprint of Training Deep Learning Models*. ICML Workshop on Challenges in Deploying and monitoring Machine Learning Systems.

Antonakakis, M., T. April, M. Bailey, M. Bernhard, E. Bursztein, J. Cochran, Z. Durumeric, et al. 2017. "Understanding the Mirai Botnet." *26th USENIX Security Symposium (USENIX Security 17)* 16: 1093–110.

Apple. 2017. *The Future Is Here: iPhone x*.

Apple. 2018. *iPhone XS and iPhone XS Max Bring the Best and Biggest Displays to iPhone*.

Apple. 2020. *Apple Unveils All-New iPad Air with A14 Bionic, Apple's Most Advanced Chip*.

Apple. 2021a. *Apple Unveils iPhone 13 Pro and iPhone 13 Pro Max*.

Apple. 2021b. *Boot Process for iPad and iPhone Devices*. Apple Platform Security.

Apple. 2021c. *Secure Software Updates*. Apple Platform Security.

Apple. 2022. *Apple Debuts iPhone 14 Pro and iPhone 14 Pro Max*.

Apple. 2023. *iPhone 15 Pro and iPhone 15 Pro Max (Product Page, Archived 20 October 2023)*.

Apple. 2024a. *Facial Matching Security*. Apple Platform Security.

Apple. 2024b. *The Secure Enclave*. Apple Platform Security.

Apple Differential Privacy Team. 2017. *Learning with Privacy at Scale*. Apple Machine Learning Research.

Arifeen, Tooba, Abdus Sami Hassan, and Jeong-A Lee. 2020. "Approximate Triple Modular Redundancy: A Survey." *IEEE Access* 8: 139851–67. https://doi.org/10.1109/access.2020.3012673.

Arivazhagan, Manoj Ghuhan, Vinay Aggarwal, Aaditya Kumar Singh, and Sunav Choudhary. 2019. "Federated Learning with Personalization Layers." *CoRR* abs/1912.00818.

Ashkboos, Saleh, Amirkeivan Mohtashami, Maximilian L. Croci, Bo Li, Pashmina Cameron, Martin Jaggi, Dan Alistarh, Torsten Hoefler, and James Hensman. 2024. "QuaRot: Outlier-Free 4-Bit Inference in Rotated LLMs." *Advances in Neural Information Processing Systems (NeurIPS)*, 100213–40. https://doi.org/10.52202/079017-3180.

Asonov, Dmitri, and Rakesh Agrawal. 2004. "Keyboard Acoustic Emanations." *IEEE Symposium on Security and Privacy, 2004. Proceedings. 2004*, 3–11. https://doi.org/10.1109/secpri.2004.1301311.

Ateniese, G., L. V. Mancini, A. Spognardi, A. Villani, D. Vitali, and G. Felici. 2015. "Hacking Smart Machines with Smarter Ones: How to Extract Meaningful Data from Machine Learning Classifiers." *International Journal of Security and Networks* 10 (3): 137. https://doi.org/10.1504/ijsn.2015.071829.

Avizienis, Algirdas, Jean-Claude Laprie, Brian Randell, and Carl Landwehr. 2004. "Basic Concepts and Taxonomy of Dependable and Secure Computing." *IEEE Transactions on Dependable and Secure Computing* 1 (1): 11–33. https://doi.org/10.1109/TDSC.2004.2.

Bai, Tao, Jinqi Luo, Jun Zhao, Bihan Wen, and Qian Wang. 2021. "Recent Advances in Adversarial Training for Adversarial Robustness." *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence*, 4312–21. https://doi.org/10.24963/ijcai.2021/591.

Bai, Yuntao et al. 2022. *Constitutional AI: Harmlessness from AI Feedback*.

Baldé, Cornelis P, Vanessa Forti, Vanessa Gray, Ruediger Kuehr, and Paul Stegmann. 2017. "The Global e-Waste Monitor 2017: Quantities, Flows and Resources." *United Nations University, International Telecommunication Union, International Solid Waste Association*.

Banbury, Colby, Vijay Janapa Reddi, Peter Torelli, Jeremy Holleman, Nat Jeffries, Csaba Kiraly, Pietro Montino, et al. 2021. "MLPerf Tiny Benchmark." *arXiv Preprint*.

Barenghi, A., G. M. Bertoni, L. Breveglieri, M. Pellicioli, and G. Pelosi. 2010. "Low Voltage Fault Attacks to AES." *2010 IEEE International Symposium on Hardware-Oriented Security and Trust (HOST)*, 7–12. https://doi.org/10.1109/hst.2010.5513121.

Barroso, Luiz André, Urs Hölzle, and Parthasarathy Ranganathan. 2019. *The Datacenter as a Computer: Designing Warehouse-Scale Machines*. Synthesis Lectures on Computer Architecture. Springer International Publishing. https://doi.org/10.1007/978-3-031-01761-2.

Basiri, A., N. Behnam, R. de Rooij, L. Hochstein, L. Kosewski, J. Reynolds, and C. Rosenthal. 2016. "Chaos Engineering." *IEEE Software* 33 (3): 35–41. https://doi.org/10.1109/ms.2016.60.

Baumann, R. 2005. "Soft Errors in Advanced Computer Systems." *IEEE Design and Test of Computers* 22 (3): 258–66. https://doi.org/10.1109/mdt.2005.69.

Baylor, Denis, Eric Breck, Heng-Tze Cheng, Noah Fiedel, Chuan Yu Foo, Zakaria Haque, Salem Haykal, et al. 2017. "TFX: A TensorFlow-Based Production-Scale Machine Learning Platform." *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1387–95. https://doi.org/10.1145/3097983.3098021.

Beaton, A. E., and J. W. Tukey. 1974. "The Fitting of Power Series, Meaning Polynomials, Illustrated on Band-Spectroscopic Data." *Technometrics* 16 (2): 147–85. https://doi.org/10.1080/00401706.1974.10489171.

Bellamy, Rachel K. E., Kuntal Dey, Michael Hind, Seung Hoffman, Sylvain Houde, Karthikeyan Kannan, Pradeep Lohia, et al. 2019. "AI Fairness 360: An Extensible Toolkit for Detecting and Mitigating Algorithmic Bias." *IBM Journal of Research and Development* 63 (4/5): 4:1–15. https://doi.org/10.1147/jrd.2019.2942287.

Bender, E. M., T. Gebru, A. McMillan-Major, and S. Shmitchell. 2021. "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?" *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, 610–23. https://doi.org/10.1145/3442188.3445922.

Bengio, Yoshua, Nicholas Léonard, and Aaron Courville. 2013. "Estimating or Propagating Gradients Through Stochastic Neurons for Conditional Computation." *arXiv Preprint arXiv:1308.3432*.

Ben-Nun, Tal, and Torsten Hoefler. 2019. "Demystifying Parallel and Distributed Deep Learning: An in-Depth Concurrency Analysis." *ACM Computing Surveys* 52 (4): 1–43. https://doi.org/10.1145/3320060.

Berger, Vance W., and YanYan Zhou. 2014. "Kolmogorov–Smirnov Test: Overview." In *Wiley StatsRef: Statistics Reference Online*. Wiley. https://doi.org/10.1002/9781118445112.stat06558.

Bernstein, Jeremy, Yu-Xiang Wang, Kamyar Azizzadenesheli, and Animashree Anandkumar. 2018. "signSGD: Compressed Optimisation for Non-Convex Problems." *International Conference on Machine Learning (ICML)* 80: 560–69.

Beyer, Betsy, Chris Jones, Jennifer Petoff, and Niall Richard Murphy. 2016. *Site Reliability Engineering: How Google Runs Production Systems*. O'Reilly Media.

Biggio, Battista, Blaine Nelson, and Pavel Laskov. 2012. "Poisoning Attacks Against Support Vector Machines." *Proceedings of the 29th International Conference on Machine Learning, ICML 2012, Edinburgh, Scotland, UK, June 26 - July 1, 2012*.

Bird, Sarah, Miro Dudík, Richard Edgar, Brandon Horn, Roman Lutz, Vanessa Milan, Mehrnoosh Sameki, Hanna Wallach, and Kathleen Walker. 2020. "Fairlearn: A Toolkit for Assessing and Improving Fairness in AI." *Microsoft Technical Report MSR-TR-2020-32*.

Birolini, A. 2017. *Reliability Engineering: Theory and Practice*. 8th ed. Springer Berlin Heidelberg. https://doi.org/10.1007/978-3-662-54209-5.

Bisong, E. 2019. "Kubeflow and Kubeflow Pipelines." Kubeflow and Kubeflow Pipelines in *Building Machine Learning and Deep Learning Models on Google Cloud Platform*. Apress. https://doi.org/10.1007/978-1-4842-4470-8_46.

Black, James R. 1969. "Electromigration–a Brief Survey and Some Recent Results." *IEEE Transactions on Electron Devices* 16 (4): 338–47. https://doi.org/10.1109/T-ED.1969.16754.

Blanchard, Peva, El Mahdi El Mhamdi, Rachid Guerraoui, and Julien Stainer. 2017. "Machine Learning with Adversaries: Byzantine Tolerant Gradient Descent." *Advances in Neural Information Processing Systems (NeurIPS)*, 119–29.

Bommasani, Rishi, Drew A. Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S. Bernstein, et al. 2021. "On the Opportunities and Risks of Foundation Models." *arXiv Preprint arXiv:2108.07258*.

Bonawitz, K., H. Eichner, W. Grieskamp, D. Huba, A. Ingerman, V. Ivanov, C. Kiddon, et al. 2019. "Towards Federated Learning at Scale: System Design." *Proceedings of Machine Learning and Systems* 3.

Bonawitz, Keith, Vladimir Ivanov, Ben Kreuter, Antonio Marcedone, H. Brendan McMahan, Sarvar Patel, Daniel Ramage, Aaron Segal, and Karn Seth. 2017. "Practical Secure Aggregation for Privacy-Preserving Machine Learning." *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 1175–91. https://doi.org/10.1145/3133956.3133982.

Bourtoule, Lucas, Varun Chandrasekaran, Christopher A. Choquette-Choo, Hengrui Jia, Adelin Travers, Baiwu Zhang, David Lie, and Nicolas Papernot. 2021. "Machine Unlearning." *2021 IEEE Symposium on Security and Privacy (SP)*, 141–59. https://doi.org/10.1109/sp40001.2021.00019.

Breier, Jakub, Xiaolu Hou, Dirmanto Jap, Lei Ma, Shivam Bhasin, and Yang Liu. 2018. "Practical Fault Attack on Deep Neural Networks." *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security* abs/1806.05859: 2204–6. https://doi.org/10.1145/3243734.3278519.

Brewer, Eric A. 2000. "Towards Robust Distributed Systems (Abstract)." *Proceedings of the Nineteenth Annual ACM Symposium on Principles of Distributed Computing*, 7. https://doi.org/10.1145/343477.343502.

Broadcom Inc. 2025. *Broadcom Ships Tomahawk 6: World's First 102.4 Tbps Switch*. Broadcom press release.

Brown, Tom B., Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. 2020. "Language Models Are Few-Shot Learners." *Advances in Neural Information Processing Systems* 33: 1877–901. https://doi.org/10.48550/arxiv.2005.14165.

Bu, Zhiqi, Jiawen Dong, Qi Long, and Weijie J. Su. 2020. "Deep Learning with Gaussian Differential Privacy." *Harvard Data Science Review* 2. https://doi.org/10.1162/99608f92.cfc5dd25.

Buolamwini, Joy, and Timnit Gebru. 2018. "Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification." *Conference on Fairness, Accountability and Transparency*, 77–91.

Burns, Brendan, Brian Grant, David Oppenheimer, Eric Brewer, and John Wilkes. 2016. "Borg, Omega, and Kubernetes." *Communications of the ACM* 59 (5): 50–57. https://doi.org/10.1145/2890784.

Bursztein, Elie, Luca Invernizzi, Karel Král, Daniel Moghimi, Jean-Michel Picod, and Marina Zhang. 2024. "Generalized Power Attacks Against Crypto Hardware Using Long-Range Deep Learning." *IACR Transactions on Cryptographic Hardware and Embedded Systems* 2024 (3): 472–99. https://doi.org/10.46586/tches.v2024.i3.472-499.

Bursztein, Elie, and Jean-Michel Picod. 2019. "A Hacker Guide to Deep Learning Based Side Channel Attacks." In *DEF CON 27*, edited by DEF CON. DEF CON.

Bushnell, Michael L, and Vishwani D Agrawal. 2002. "Built-in Self-Test." *Essentials of Electronic Testing for Digital, Memory and Mixed-Signal VLSI Circuits* 1: 489–548. https://doi.org/10.1007/0-306-47040-3_15.

Buyya, Rajkumar, Anton Beloglazov, and Jemal Abawajy. 2010. "Energy-Efficient Management of Data Center Resources for Cloud Computing: A Vision, Architectural Elements, and Open Challenges." *arXiv Preprint arXiv:1006.0308*, 6–17. https://doi.org/10.48550/arXiv.1006.0308.

Cai, Han, Chuang Gan, Ligeng Zhu, and Song Han. 2020. "TinyTL: Reduce Activations, Not Trainable Parameters for Efficient on-Device Learning." *Advances in Neural Information Processing Systems* 33: 11285–97.

Cai, Tianle, Yuhong Li, Zhengyang Geng, Hongwu Peng, Jason D. Lee, Deming Chen, and Tri Dao. 2024. "Medusa: Simple LLM Inference Acceleration Framework with Multiple Decoding Heads." *Proceedings of the 41st International Conference on Machine Learning (ICML)*.

Calders, Toon, and Sicco Verwer. 2010. "Three Naive Bayes Approaches for Discrimination-Free Classification." *Data Mining and Knowledge Discovery* 21 (2): 277–92. https://doi.org/10.1007/s10618-010-0190-x.

California Legislature. 2023. *California Consumer Privacy Act of 2018*. California Legislative Information.

Cao, Yinzhi, and Junfeng Yang. 2015. "Towards Making Systems Forget with Machine Unlearning." *2015 IEEE Symposium on Security and Privacy*, 463–80. https://doi.org/10.1109/sp.2015.35.

Carlini, Nicholas, Jamie Hayes, Milad Nasr, Matthew Jagielski, Vikash Sehwag, Florian Tramèr, Borja Balle, Daphne Ippolito, and Eric Wallace. 2023. "Extracting Training Data from Diffusion Models." *32nd USENIX Security Symposium (USENIX Security 23)*, 5253–70.

Carlini, Nicholas, and David Wagner. 2017. "Towards Evaluating the Robustness of Neural Networks." *2017 IEEE Symposium on Security and Privacy (SP)*, 39–57. https://doi.org/10.1109/sp.2017.49.

Carlini, N., D. Paleka, K. D. Dvijotham, T. Steinke, J. Hayase, A. F. Cooper, K. Lee, et al. 2024. "Stealing Part of a Production Language Model." *arXiv Preprint arXiv:2403.06634*.

Carlini, N., F. Tramer, E. Wallace, M. Jagielski, A. Herbert-Voss, K. Lee, A. Roberts, et al. 2021. "Extracting Training Data from Large Language Models." *30th USENIX Security Symposium (USENIX Security 21)*, 2633–50.

Caveness, E., P. S. G. C., Z. Peng, N. Polyzotis, S. Roy, and M. Zinkevich. 2020. "TensorFlow Data Validation: Data Analysis and Validation in Continuous ML Pipelines." *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*, 2793–96. https://doi.org/10.1145/3318464.3384707.

Cavoukian, A. 2012. "Privacy by Design: Origins, Meaning, and Prospects for Assuring Privacy and Trust in the Information Era." *Office of the Information and Privacy Commissioner* 1: 170–208. https://doi.org/10.4018/978-1-61350-501-4.ch007.

Cenci, Marcelo Pilotto, Tatiana Scarazzato, Daniel Dotto Munchen, Paula Cristina Dartora, Hugo Marcelo Veit, Andrea Moura Bernardes, and Pablo R. Dias. 2021. "Eco-Friendly Electronics—a Comprehensive Review." *Advanced Materials Technologies* 7 (2): 2001263. https://doi.org/10.1002/admt.202001263.

Center, Pew Research. 2023. *What Americans Know about AI, Cybersecurity and Big Tech*.

Centre for International Governance Innovation. 2023. *Who Is Responsible When Autonomous Systems Fail?* CIGI policy article.

Cerebras Systems. 2024. *Cerebras Systems, Inc. Registration Statement on Form S-1*.

Chandola, Varun, Arindam Banerjee, and Vipin Kumar. 2009. "Anomaly Detection: A Survey." *ACM Computing Surveys* 41 (3): 1–58. https://doi.org/10.1145/1541880.1541882.

Chapelle, Olivier, Thorsten Joachims, Filip Radlinski, and Yisong Yue. 2012. "Large-Scale Validation and Analysis of Interleaved Search Evaluation." *ACM Transactions on Information Systems* 30 (1): 1–41. https://doi.org/10.1145/2094072.2094078.

Chen, C., S. Borgeaud, G. Irving, J.-B. Lespiau, L. Sifre, and J. Jumper. 2023. "Accelerating Large Language Model Decoding with Speculative Sampling." *arXiv Preprint arXiv:2302.01318*.

Chen, Chaofan, Oscar Li, Daniel Tao, Alina Barnett, Cynthia Rudin, and Jonathan Su. 2019. "This Looks Like That: Deep Learning for Interpretable Image Recognition." In *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, edited by Hanna M. Wallach, Hugo Larochelle, Alina Beygelzimer, Florence d'Alché-Buc, Emily B. Fox, and Roman Garnett. Curran Associates.

Chen, H.-W. 2006. "Gallium, Indium, and Arsenic Pollution of Groundwater from a Semiconductor Manufacturing Area of Taiwan." *Bulletin of Environmental Contamination and Toxicology* 77 (2): 289–96. https://doi.org/10.1007/s00128-006-1062-3.

Chen, Liyuan, Zihao Ye, Yonghao Wu, Danyang Zhuo, Luis Ceze, and Arvind Krishnamurthy. 2024. "Punica: Multi-Tenant LoRA Serving." *Proceedings of the 7th Conference on Machine Learning and Systems (MLSys)*.

Chen, Tianqi, Bing Xu, Chiyuan Zhang, and Carlos Guestrin. 2016. "Training Deep Nets with Sublinear Memory Cost." *arXiv Preprint arXiv:1604.06174*.

Chen, Ting, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. 2020. "A Simple Framework for Contrastive Learning of Visual Representations." *Proceedings of the 37th International Conference on Machine Learning*, ICML'20, vol. 119: 1597–607.

Chen, Zitao, Niranjhana Narayanan, Bo Fang, Guanpeng Li, Karthik Pattabiraman, and Nathan DeBardeleben. 2020. "TensorFI: A Flexible Fault Injection Framework for TensorFlow Applications." *2020 IEEE 31st International Symposium on Software Reliability Engineering (ISSRE)*, 426–35. https://doi.org/10.1109/issre5003.2020.00047.

Choquette, Jack. 2023. "NVIDIA Hopper H100 GPU: Scaling Performance." *IEEE Micro* 43 (3): 9–17. https://doi.org/10.1109/mm.2023.3256796.

Choquette, Jack, Wishwesh Gandhi, Olivier Giroux, Nick Stam, and Ronny Krashinsky. 2021. "NVIDIA A100 Tensor Core GPU: Performance and Innovation." *IEEE Micro* 41 (2): 29–35. https://doi.org/10.1109/mm.2021.3061394.

Choquette, Jack, Olivier Giroux, and Denis Foley. 2018. "Volta: Performance and Programmability." *IEEE Micro* 38 (2): 42–52. https://doi.org/10.1109/MM.2018.022071134.

Chouldechova, Alexandra. 2017. "Fair Prediction with Disparate Impact: A Study of Bias in Recidivism Prediction Instruments." *Big Data* 5 (2): 153–63. https://doi.org/10.1089/big.2016.0047.

Chowdhery, Aakanksha, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, et al. 2022. "PaLM: Scaling Language Modeling with Pathways." *arXiv Preprint arXiv:2204.02311*.

Christiano, P. F., J. Leike, T. B. Brown, M. Martic, S. Legg, and D. Amodei. 2017. "Deep Reinforcement Learning from Human Preferences." In *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, edited by Isabelle Guyon, Ulrike von Luxburg, Samy Bengio, Hanna M. Wallach, Rob Fergus, S. V. N. Vishwanathan, and Roman Garnett. Curran Associates.

Chung, Jae-Won, Yile Gu, Insu Jang, Luoxi Meng, Nikhil Bansal, and Mosharaf Chowdhury. 2023. "Reducing Energy Bloat in Large Model Training." *Proceedings of the ACM SIGOPS 30th Symposium on Operating Systems Principles* abs/2312.06902: 144–59. https://doi.org/10.1145/3694715.3695970.

Clos, C. 1953. "A Study of Non-Blocking Switching Networks." *Bell System Technical Journal* 32 (2): 406–24. https://doi.org/10.1002/j.1538-7305.1953.tb01433.x.

Cloud Native Computing Foundation. 2024. *Kubernetes: Production-Grade Container Orchestration*.

Cloudflare. 2019. *Details of the Cloudflare Outage on July 2, 2019*. Cloudflare incident report.

CNBC. 2023. *California DMV Suspends Cruise's Self-Driving Car Permits*. CNBC news article.

Coffman, Edward G., Michael Elphick, and Arie Shoshani. 1971. "System Deadlocks." *ACM Computing Surveys* 3 (2): 67–78. https://doi.org/10.1145/356586.356588.

Cohen, Jeremy, Elan Rosenfeld, and Zico Kolter. 2019. "Certified Adversarial Robustness via Randomized Smoothing." *International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 97: 1310–20.

Constantinescu, C. 2008. "Intermittent Faults and Effects on Reliability of Integrated Circuits." *2008 Annual Reliability and Maintainability Symposium*, 370–74. https://doi.org/10.1109/rams.2008.4925824.

Cooper, Tom, Suzanne Fallender, Joyann Pafumi, Jon Dettling, Sebastien Humbert, and Lindsay Lessard. 2011. "A Semiconductor Company's Examination of Its Water Footprint Approach." *Proceedings of the 2011 IEEE International Symposium on Sustainable Systems and Technology*, 1–6. https://doi.org/10.1109/issst.2011.5936865.

Cottier, Ben, Robi Rahman, Loredana Fattorini, Nestor Maslej, Tamay Besiroglu, and David Owen. 2024. "The Rising Costs of Training Frontier AI Models." *arXiv Preprint arXiv:2405.21015*.

Covington, Paul, Jay Adams, and Emre Sargin. 2016. "Deep Neural Networks for YouTube Recommendations." *Proceedings of the 10th ACM Conference on Recommender Systems*, 191–98. https://doi.org/10.1145/2959100.2959190.

Croce, Francesco, Maksym Andriushchenko, Vikash Sehwag, Edoardo Debenedetti, Nicolas Flammarion, Mung Chiang, Prateek Mittal, and Matthias Hein. 2020. "RobustBench: A Standardized Adversarial Robustness Benchmark." *arXiv Preprint arXiv:2010.09670*.

Croce, Francesco, and Matthias Hein. 2020. "Reliable Evaluation of Adversarial Robustness with an Ensemble of Diverse Parameter-Free Attacks." *arXiv Preprint arXiv:2003.01690*.

Culler, David, Richard Karp, David Patterson, Abhijit Sahay, Klaus Erik Schauser, Eunice Santos, Ramesh Subramonian, and Thorsten von Eicken. 1993. "LogP: Towards a Realistic Model of Parallel Computation." *ACM SIGPLAN Notices* 28 (7): 1–12. https://doi.org/10.1145/173284.155333.

Cyberspace Administration of China, National Development and Reform Commission, Ministry of Education, Ministry of Science and Technology, Ministry of Industry and Information Technology, Ministry of Public Security, and National Radio and Television Administration. 2023. *Interim Measures for Administration of Generative Artificial Intelligence Services*. State Council Gazette.

Daly, J. T. 2006. "A Higher Order Estimate of the Optimum Checkpoint Interval for Restart Dumps." *Future Generation Computer Systems* 22 (3): 303–12. https://doi.org/10.1016/j.future.2004.11.016.

Dao, T., D. Y. Fu, S. Ermon, A. Rudra, and C. Ré. 2022. "FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness." *Advances in Neural Information Processing Systems (NeurIPS)* 35: 16344–59. https://doi.org/10.52202/068431-1189.

Dastin, Jeffrey. 2018. "Amazon Scraps Secret AI Recruiting Tool That Showed Bias Against Women." *Reuters*.

Databricks. 2026. *Mosaic Streaming: Main Concepts*. Databricks Mosaic AI Training Documentation.

Dayarathna, Miyuru, Yonggang Wen, and Rui Fan. 2016. "Data Center Energy Consumption Modeling: A Survey." *IEEE Communications Surveys & Tutorials* 18 (1): 732–94. https://doi.org/10.1109/comst.2015.2481183.

De, Soham, Leonard Berrada, Jamie Hayes, Samuel L. Smith, and Borja Balle. 2022. "Unlocking High-Accuracy Differentially Private Image Classification Through Scale." *arXiv Preprint arXiv:2204.13650*.

Dean, Jeff. 2024. *Exciting Directions in Systems for Machine Learning*. MLSys 2024 invited talk.

Dean, Jeffrey, and Luiz André Barroso. 2013a. "The Tail at Scale." *Communications of the ACM* 56 (2): 74–80. https://doi.org/10.1145/2408776.2408794.

Dean, Jeffrey, and Luiz André Barroso. 2013b. "The Tail at Scale." *Communications of the ACM* 56 (2): 74–80.

Dean, Jeffrey, Greg Corrado, Rajat Monga, Kai Chen, Matthieu Devin, Quoc V. Le, Mark Z. Mao, et al. 2012. "Large Scale Distributed Deep Networks." In *Advances in Neural Information Processing Systems (NeurIPS)*, edited by Peter L. Bartlett, Fernando C. N. Pereira, Christopher J. C. Burges, Léon Bottou, and Kilian Q. Weinberger, vol. 25. Curran Associates.

Dean, Jeffrey, and Sanjay Ghemawat. 2004. "MapReduce: Simplified Data Processing on Large Clusters." *Proceedings of the 6th Symposium on Operating Systems Design and Implementation (OSDI)*, 137–50.

DeepSeek-AI. 2025. "Insights into DeepSeek-V3: Scaling Challenges and Reflections on Hardware for AI Architectures." *Proceedings of the 52nd Annual International Symposium on Computer Architecture*, 1731–45. https://doi.org/10.1145/3695053.3731412.

DeepSpeed Developers. 2026a. *DeepSpeed Model Checkpointing*. DeepSpeed documentation.

DeepSpeed Developers. 2026b. *DeepSpeed Universal Checkpointing with DeepSpeed: A Practical Guide*. DeepSpeed tutorial.

Dell Technologies. 2026. *Data Center Power and Cooling Solutions*.

Dennard, Robert H., Frank H. Gaensslen, Hwa-Nien Yu, Victor L. Rideout, Elias Bassous, and Antoine R. LeBlanc. 1974. "Design of Ion-Implanted MOSFET's with Very Small Physical Dimensions." *IEEE J. Solid-State Circuits* 9 (5): 256–68. https://doi.org/10.1109/jssc.1974.1050511.

Department for Education, and Ofqual. 2020. *GCSE and A Level Students to Receive Centre Assessment Grades*. GOV.UK News Story.

Dettmers, Tim, Mike Lewis, Younes Belkada, and Luke Zettlemoyer. 2022. "LLM.int8(): 8-Bit Matrix Multiplication for Transformers at Scale." *Advances in Neural Information Processing Systems (NeurIPS)*, 30318–32. https://doi.org/10.52202/068431-2198.

Devarajan, Hariharan, Huihuo Zheng, Anthony Kougkas, Xian-He Sun, and Venkatram Vishwanath. 2021. "DLIO: A Data-Centric Benchmark for Scientific Deep Learning Applications." *2021 IEEE/ACM 21st International Symposium on Cluster, Cloud and Internet Computing (CCGrid)*, 81–91. https://doi.org/10.1109/CCGrid51090.2021.00018.

Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. "BERT: Pre-Training of Deep Bidirectional Transformers for Language Understanding." *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 4171–86. https://doi.org/10.18653/v1/n19-1423.

Dixit, H. D., S. Pendharkar, M. Beadon, C. Mason, T. Chakravarthy, B. Muthiah, and S. Sankar. 2021. *Silent Data Corruptions at Scale*. arXiv preprint arXiv:2102.11245.

Dodge, Jesse, Taylor Prewitt, Remi Tachet des Combes, Erika Odmark, Roy Schwartz, Emma Strubell, Alexandra Sasha Luccioni, Noah A. Smith, Nicole DeCario, and Will Buchanan. 2022. "Measuring the Carbon Intensity of AI in Cloud Instances." *2022 ACM Conference on Fairness Accountability and Transparency*, 1877–94. https://doi.org/10.1145/3531146.3533234.

Du, Nan, Yanping Huang, Andrew M. Dai, Simon Tong, Dmitry Lepikhin, Yuanzhong Xu, Maxim Krikun, et al. 2022. "GLaM: Efficient Scaling of Language Models with Mixture-of-Experts." In *International Conference on Machine Learning (ICML)*, edited by Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu, and Sivan Sabato, vol. 162. Proceedings of Machine Learning Research. PMLR.

Dubey, Abhimanyu, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, et al. 2024. *The Llama 3 Herd of Models*. arXiv preprint arXiv:2407.21783.

Dutta, Sanghamitra, Gauri Joshi, Soumyadip Ghosh, Parijat Dube, and Priya Nagpurkar. 2018. "Slow and Stale Gradients Can Win the Race: Error-Runtime Trade-Offs in Distributed SGD." *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Statistics* 84: 803–12.

Dwork, Cynthia, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Richard Zemel. 2012. "Fairness Through Awareness." *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, 214–26. https://doi.org/10.1145/2090236.2090255.

Dwork, Cynthia, Frank McSherry, Kobbi Nissim, and Adam Smith. 2006. "Calibrating Noise to Sensitivity in Private Data Analysis." In *Theory of Cryptography Conference (TCC)*, edited by Shai Halevi and Tal Rabin, vol. 3876. Lecture Notes in Computer Science. Springer Berlin Heidelberg. https://doi.org/10.1007/11681878_14.

Dwork, Cynthia, and Aaron Roth. 2014. "The Algorithmic Foundations of Differential Privacy." *Foundations and Trends® in Theoretical Computer Science*, Foundations and trends in theoretical computer science, vol. 9 (3-4): 211–487. https://doi.org/10.1561/0400000042.

Ebrahimi, K., G. F. Jones, and A. S. Fleischer. 2014. "A Review of Data Center Cooling Technology, Operating Conditions and the Corresponding Low-Grade Waste Heat Recovery Opportunities." *Renewable and Sustainable Energy Reviews* 31: 622–38. https://doi.org/10.1016/j.rser.2013.12.007.

Eckles, Dean, Brian Karrer, and Johan Ugander. 2017. "Design and Analysis of Experiments in Networks: Reducing Bias from Interference." *Journal of Causal Inference* 5 (1): 1–23. https://doi.org/10.1515/jci-2015-0021.

Elnozahy, Elmootazbellah N., Lorenzo Alvisi, Yi-Min Wang, and David B. Johnson. 2002. "A Survey of Rollback-Recovery Protocols in Message-Passing Systems." *ACM Computing Surveys* 34 (3): 375–408. https://doi.org/10.1145/568522.568525.

Elsken, Thomas, Jan Hendrik Metzen, and Frank Hutter. 2019. "Neural Architecture Search." In *The Springer Series on Challenges in Machine Learning*, vol. 20. Springer International Publishing. https://doi.org/10.1007/978-3-030-05318-5_3.

Elsworth, Cooper, Keguo Huang, David Patterson, Ian Schneider, Robert Sedivy, Savannah Goodman, Ben Townsend, et al. 2025. "Measuring the Environmental Impact of Delivering AI at Google Scale." *arXiv Preprint arXiv:2508.15734*.

Emer, Joel S., and Douglas W. Clark. 1984. "A Characterization of Processor Performance in the VAX-11/780." *ACM SIGARCH Computer Architecture News* 12 (3): 301–10. https://doi.org/10.1145/773453.808199.

Engineering, M. 2016. *Introducing FBLearner Flow: Facebook's AI Backbone*. Engineering at Meta Blog.

Engineering, S. 2019. *The Winding Road to Better Machine Learning Infrastructure Through TensorFlow Extended and Kubeflow*. Spotify Engineering Blog.

Environment and Climate Change Canada. 2026. *Emission Factors and Reference Values*.

Environmental Management—Water Footprint—Principles, Requirements and Guidelines, Pub. L. Nos. ISO 14046:2014 (2014).

Epoch AI. 2026a. *Data on AI Supercomputers*.

Epoch AI. 2026b. *Machine Learning Hardware*.

Epoch AI. 2026c. *Notable AI Models*.

Erlang, Agner Krarup. 1909. "The Theory of Probabilities and Telephone Conversations." *Nyt Tidsskrift for Matematik B* 20: 33–39.

ETH Zurich Computer Vision Lab. 2026. *AI Benchmark: Mobile SoC Ranking*.

Ethernet Alliance. 2025. *2025 Ethernet Roadmap*. Ethernet Alliance roadmap.

European Data Protection Board. 2018. *Guidelines on Automated Individual Decision-Making and Profiling for the Purposes of Regulation 2016/679*.

European Parliament and Council of the European Union. 2024. *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying down Harmonised Rules on Artificial Intelligence (AI Act)*. Official Journal of the European Union, L 2024/1689.

European Parliament, and Council of the European Union. 2016. *Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016*. Official Journal of the European Union.

Evans, Richard, and Jim Gao. 2016. *DeepMind AI Reduces Google Data Centre Cooling Bill by 40%*. DeepMind Blog.

Executive Office of the President. 2023. *Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence*. Executive Order 14110, 88 FR 75191.

Executive Office of the President. 2025a. *Initial Rescissions of Harmful Executive Orders and Actions*. Executive Order 14148, 90 FR 8237.

Executive Office of the President. 2025b. *Removing Barriers to American Leadership in Artificial Intelligence*. Executive Order 14179.

Executive Office of the President. 2026. *Promoting Advanced Artificial Intelligence Innovation and Security*. Executive Order 14409, 91 FR 34565.

Eykholt, Kevin, Ivan Evtimov, Earlence Fernandes, Bo Li, Amir Rahmati, Chaowei Xiao, Atul Prakash, Tadayoshi Kohno, and Dawn Song. 2017. "Robust Physical-World Attacks on Deep Learning Models." *ArXiv Preprint* abs/1707.08945 (July).

Eykholt, K., I. Evtimov, E. Fernandes, B. Li, A. Rahmati, C. Xiao, A. Prakash, T. Kohno, and D. Song. 2018. "Robust Physical-World Attacks on Deep Learning Visual Classification." *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition* 1707.08945: 1625–34. https://doi.org/10.1109/cvpr.2018.00175.

Farwell, J. P., and R. Rohozinski. 2011. "Stuxnet and the Future of Cyber War." *Survival* 53 (1): 23–40. https://doi.org/10.1080/00396338.2011.555586.

Federal Aviation Administration. 2015. *Airworthiness Directive: Boeing 787 Generator Control Units Failsafe Reset*. FAA Airworthiness Directive 2015-10066.

Federal Trade Commission. 2021. *Nixing the Fix: An FTC Report to Congress on Repair Restrictions*. Federal Trade Commission.

Fedus, William, Barret Zoph, and Noam Shazeer. 2022. "Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity." *Journal of Machine Learning Research* 23 (120): 1–39.

Feldman, M., S. A. Friedler, J. Moeller, C. Scheidegger, and S. Venkatasubramanian. 2015. "Certifying and Removing Disparate Impact." *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 259–68. https://doi.org/10.1145/2783258.2783311.

Finn, Chelsea, Pieter Abbeel, and Sergey Levine. 2017. "Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks." *Proceedings of the 34th International Conference on Machine Learning (ICML)*.

Foley, Denis, and John Danskin. 2017. "Ultra-Performance Pascal GPU and NVLink Interconnect." *IEEE Micro* 37 (2): 7–17. https://doi.org/10.1109/MM.2017.37.

Forti, V., C. P. Baldé, R. Kuehr, and G. Bel. 2020. *The Global e-Waste Monitor 2020: Quantities, Flows and the Circular Economy Potential*. United Nations University/United Nations Institute for Training and Research; International Telecommunication Union; International Solid Waste Association.

Frantar, Elias, Saleh Ashkboos, Torsten Hoefler, and Dan Alistarh. 2023. "GPTQ: Accurate Post-Training Quantization for Generative Pre-Trained Transformers." In *CoRR*, abs/2210.17323. https://doi.org/10.48550/ARXIV.2210.17323.

Fredrikson, Matt, Somesh Jha, and Thomas Ristenpart. 2015. "Model Inversion Attacks That Exploit Confidence Information and Basic Countermeasures." *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, 1322–33. https://doi.org/10.1145/2810103.2813677.

Friedman, B. 1996. "Value-Sensitive Design." *Interactions* 3 (6): 16–23. https://doi.org/10.1145/242485.242493.

Gal, Yarin, and Zoubin Ghahramani. 2016. "Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning." *Proceedings of the 33rd International Conference on Machine Learning (ICML)* 48: 1050–59.

Gama, João, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. 2014. "A Survey on Concept Drift Adaptation." *ACM Computing Surveys* 46 (4): 1–37. https://doi.org/10.1145/2523813.

Gandolfi, Karine, Christophe Mourtel, and Francis Olivier. 2001. "Electromagnetic Analysis: Concrete Results." In *Cryptographic Hardware and Embedded Systems — CHES 2001*. Springer Berlin Heidelberg. https://doi.org/10.1007/3-540-44709-1_21.

Gangidi, Adi, Rui Miao, Sandeep Hebbani, Gaya Nagarajan, Omar Baldonado, Lixin Gao, Hany Morsy Goes, et al. 2024. "RDMA over Ethernet for Distributed AI Training at Meta Scale." *Proceedings of the ACM SIGCOMM 2024 Conference*, 56–69. https://doi.org/10.1145/3651890.3672233.

Gao, J. 2014. *Machine Learning Applications for Data Center Optimization*. Google; Google White Paper.

Gao, Y., S. F. Al-Sarawi, and D. Abbott. 2020. "Physical Unclonable Functions." *Nature Electronics* 3 (2): 81–91. https://doi.org/10.1038/s41928-020-0372-5.

Gassend, Blaise, Dwaine Clarke, Marten van Dijk, and Srinivas Devadas. 2002. "Silicon Physical Random Functions." *Proceedings of the 9th ACM Conference on Computer and Communications Security*, 148–60. https://doi.org/10.1145/586110.586132.

Gebru, Timnit, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. 2021. "Datasheets for Datasets." *Communications of the ACM* 64 (12): 86–92. https://doi.org/10.1145/3458723.

Geiger, Atticus, Hanson Lu, Thomas Icard, and Christopher Potts. 2021. "Causal Abstractions of Neural Networks." In *Advances in Neural Information Processing Systems 34: Annual Conference on Neural Information Processing Systems 2021, NeurIPS 2021, December 6-14, 2021, Virtual*, edited by Marc'Aurelio Ranzato, Alina Beygelzimer, Yann N. Dauphin, Percy Liang, and Jennifer Wortman Vaughan. Curran Associates.

Genkin, Daniel, Adi Shamir, and Eran Tromer. 2017. "Acoustic Cryptanalysis." *Journal of Cryptology* 30: 392–443. https://doi.org/10.1007/s00145-015-9224-2.

Gentry, C. 2009. "Fully Homomorphic Encryption Using Ideal Lattices." *Proceedings of the Forty-First Annual ACM Symposium on Theory of Computing*, 169–78. https://doi.org/10.1145/1536414.1536440.

Ghodsi, Ali, Matei Zaharia, Benjamin Hindman, Andy Konwinski, Scott Shenker, and Ion Stoica. 2011. "Dominant Resource Fairness: Fair Allocation of Multiple Resource Types." *8th USENIX Symposium on Networked Systems Design and Implementation (NSDI 11)*, 323–36.

Gholami, Amir, Sehoon Kim, Zhen Dong, Zhewei Yao, Michael W. Mahoney, and Kurt Keutzer. 2022. "A Survey of Quantization Methods for Efficient Neural Network Inference." In *Low-Power Computer Vision: Improve the Efficiency of Artificial Intelligence*. Chapman; Hall/CRC. https://doi.org/10.1201/9781003162810-13.

Gibiansky, A. 2017. *Bringing HPC Techniques to Deep Learning*. Baidu Research Technical Blog.

Gilbert, Seth, and Nancy Lynch. 2002. "Brewer's Conjecture and the Feasibility of Consistent, Available, Partition-Tolerant Web Services." *ACM SIGACT News* 33 (2): 51–59. https://doi.org/10.1145/564585.564601.

GitLab. 2017. *Postmortem of Database Outage of January 31*. GitLab postmortem.

Gnad, D. R. E., F. Oboril, and M. B. Tahoori. 2017. "Voltage Drop-Based Fault Attacks on FPGAs Using Valid Bitstreams." *2017 27th International Conference on Field Programmable Logic and Applications (FPL)*, 1–7. https://doi.org/10.23919/fpl.2017.8056840.

Goldberg, David. 1991. "What Every Computer Scientist Should Know about Floating-Point Arithmetic." *ACM Computing Surveys* 23 (1): 5–48. https://doi.org/10.1145/103162.103163.

Gomez-Uribe, Carlos A., and Neil Hunt. 2015. "The Netflix Recommender System: Algorithms, Business Value, and Innovation." *ACM Transactions on Management Information Systems* 6 (4): 1–19. https://doi.org/10.1145/2843948.

Goncalves, A., P. Ray, B. Soper, J. Stevens, L. Coyle, and A. P. Sales. 2020. "Generation and Evaluation of Synthetic Patient Data." *BMC Medical Research Methodology* 20 (1): 1–40. https://doi.org/10.1186/s12874-020-00977-1.

Goodfellow, I. J., J. Shlens, and C. Szegedy. 2015. "Explaining and Harnessing Adversarial Examples." *International Conference on Learning Representations (ICLR)*.

Goodfellow, Ian, Yoshua Bengio, and Aaron Courville. 2016. *Deep Learning*. MIT Press.

Goodhart, Charles A. E. 1984. "Problems of Monetary Management: The UK Experience." In *Monetary Theory and Practice*. Palgrave. https://doi.org/10.1007/978-1-349-17295-5_4.

Google. 2023. *Power Usage Effectiveness: Google Data Centers*.

Google. 2024. *ShieldGemma: Generative AI Content Moderation Based on Gemma*. Google AI for Developers documentation.

Google Cloud. 2026a. *About Cloud Storage Objects*. Google Cloud Documentation.

Google Cloud. 2026b. *TPU v6e*.

Government of Ireland. 2022. *Government Statement on the Role of Data Centres in Ireland's Enterprise Strategy*. Government policy statement.

Goyal, Priya, Piotr Dollár, Ross Girshick, Pieter Noordhuis, Lukasz Wesolowski, Aapo Kyrola, Andrew Tulloch, Yangqing Jia, and Kaiming He. 2017. "Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour." *arXiv Preprint arXiv:1706.02677* abs/1706.02677.

Gräfe, Ralf, Qutub Syed Sha, Florian Geissler, and Michael Paulitsch. 2023. "Large-Scale Application of Fault Injection into PyTorch Models -an Extension to PyTorchFI for Validation Efficiency." *2023 53rd Annual IEEE/IFIP International Conference on Dependable Systems and Networks - Supplemental Volume (DSN-s)*, 56–62. https://doi.org/10.1109/dsn-s58398.2023.00025.

Graham, R. L., D. Bureddy, P. Lui, H. Rober, G. Bloch, G. Shainer, J. Poole, et al. 2020. "Scalable Hierarchical Aggregation and Reduction Protocol (SHARP) Streaming-Aggregation Hardware Design and Evaluation." *International Conference on High Performance Computing*, 41–59. https://doi.org/10.1007/978-3-030-50743-5_3.

Gray, J. N. 1978. "Notes on Data Base Operating Systems." In *Operating Systems: An Advanced Course*. Springer Berlin Heidelberg. https://doi.org/10.1007/3-540-08755-9_9.

Gray, Jim. 1986. "Why Do Computers Stop and What Can Be Done about It?" *Proceedings of the 5th Symposium on Reliability in Distributed Software and Database Systems*, 3–12.

Greshake, Kai, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. 2023. "Not What You've Signed up for: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection." *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*, 79–90. https://doi.org/10.1145/3605764.3623985.

Grossman, E. 2007. *High Tech Trash: Digital Devices, Hidden Toxics, and Human Health*. Island press.

Gu, J., M. Chowdhury, K. G. Shin, Y. Zhu, M. Jeon, J. Qian, H. H. Liu, and C. Guo. 2019. "Tiresias: A GPU Cluster Manager for Distributed Deep Learning." *Proceedings of the 16th USENIX Symposium on Networked Systems Design and Implementation (NSDI '19)*, 485–500.

Gu, Tianyu, Brendan Dolan-Gavitt, and Siddharth Garg. 2017. "BadNets: Identifying Vulnerabilities in the Machine Learning Model Supply Chain." *arXiv Preprint arXiv:1708.06733*.

Guo, Chuan, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. "On Calibration of Modern Neural Networks." *International Conference on Machine Learning (ICML)* 70: 1321–30.

Guo, Chuanxiong, Haitao Wu, Zhong Deng, Gaurav Soni, Jianxi Ye, Jitendra Padhye, and Marina Lipshteyn. 2016. "RDMA over Commodity Ethernet at Scale." *Proceedings of the 2016 ACM SIGCOMM Conference*, 202–15. https://doi.org/10.1145/2934872.2934908.

Gupta, Maanak, Charankumar Akiri, Kshitiz Aryal, Eli Parker, and Lopamudra Praharaj. 2023. "From ChatGPT to ThreatGPT: Impact of Generative AI in Cybersecurity and Privacy." *IEEE Access* 11: 80218–45. https://doi.org/10.1109/access.2023.3300381.

Gupta, Udit, Young Geun Kim, Sylvia Lee, Jordan Tse, Hsien-Hsin S Lee, Gu-Yeon Wei, David Brooks, and Carole-Jean Wu. 2022. "Chasing Carbon: The Elusive Environmental Footprint of Computing." *IEEE Micro* 42 (4): 37–47. https://doi.org/10.1109/MM.2022.3163226.

Gupta, Udit, Carole-Jean Wu, Xiaodong Wang, Maxim Naumov, Brandon Reagen, David Brooks, Bradford Cottel, et al. 2020. "The Architectural Implications of Facebook's DNN-Based Personalized Recommendation." *2020 IEEE International Symposium on High Performance Computer Architecture (HPCA)*, 488–501. https://doi.org/10.1109/HPCA47549.2020.00047.

Gustafson, John L. 1988. "Reevaluating Amdahl's Law." *Communications of the ACM* 31 (5): 532–33. https://doi.org/10.1145/42411.42415.

Hamming, R. W. 1950. "Error Detecting and Error Correcting Codes." *Bell System Technical Journal* 29 (2): 147–60. https://doi.org/10.1002/j.1538-7305.1950.tb00463.x.

Han, Song, Jeff Pool, John Tran, and William J. Dally. 2015. "Learning Both Weights and Connections for Efficient Neural Networks." *Advances in Neural Information Processing Systems 28 (NeurIPS 2015)*, 1135–43.

Hard, Andrew, Kanishka Rao, Rajiv Mathews, Swaroop Ramaswamy, Françoise Beaufays, Sean Augenstein, Hubert Eichner, Chloé Kiddon, and Daniel Ramage. 2018. "Federated Learning for Mobile Keyboard Prediction." *arXiv Preprint arXiv:1811.03604*.

Hardt, Moritz, Eric Price, and Nathan Srebro. 2016. "Equality of Opportunity in Supervised Learning." *Advances in Neural Information Processing Systems (NeurIPS)* 29: 3315–23.

Hayes, T. L., K. Kafle, R. Shrestha, M. Acharya, and C. Kanan. 2020. "REMIND Your Neural Network to Prevent Catastrophic Forgetting." In *Computer Vision – ECCV 2020*. Springer. https://doi.org/10.1007/978-3-030-58598-3_28.

Hazelwood, Kim, Sarah Bird, David Brooks, Soumith Chintala, Utku Diril, Dmytro Dzhulgakov, Mohamed Fawzy, et al. 2018. "Applied Machine Learning at Facebook: A Datacenter Infrastructure Perspective." *2018 IEEE International Symposium on High Performance Computer Architecture (HPCA)*, 620–29. https://doi.org/10.1109/hpca.2018.00059.

He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. "Deep Residual Learning for Image Recognition." *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770–78. https://doi.org/10.1109/cvpr.2016.90.

He, K., X. Chen, S. Xie, Y. Li, P. Dollár, and R. Girshick. 2021. "Masked Autoencoders Are Scalable Vision Learners." *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* 18: 15979–88. https://doi.org/10.1109/cvpr52688.2022.01553.

He, Yi, Prasanna Balaprakash, and Yanjing Li. 2020. "FIdelity: Efficient Resilience Analysis Framework for Deep Learning Accelerators." *2020 53rd Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*, 270–81. https://doi.org/10.1109/micro50266.2020.00033.

He, Yi, Mike Hutton, Steven Chan, Robert De Gruijl, Rama Govindaraju, Nishant Patil, and Yanjing Li. 2023. "Understanding and Mitigating Hardware Failures in Deep Learning Training Systems." *Proceedings of the 50th Annual International Symposium on Computer Architecture*, 1–16. https://doi.org/10.1145/3579371.3589105.

Hébert-Johnson, U., M. P. Kim, O. Reingold, and G. N. Rothblum. 2018. "Multicalibration: Calibration for the (Computationally-Identifiable) Masses." In *Proceedings of the 35th International Conference on Machine Learning, ICML 2018, Stockholmsmässan, Stockholm, Sweden, July 10-15, 2018*, edited by Jennifer G. Dy and Andreas Krause, vol. 80. Proceedings of Machine Learning Research. PMLR.

Henderson, Peter, Jieru Hu, Joshua Romoff, Emma Brunskill, Dan Jurafsky, and Joelle Pineau. 2020. "Towards the Systematic Reporting of the Energy and Carbon Footprints of Machine Learning." *CoRR* abs/2002.05651 (248): 1–43. https://doi.org/10.48550/arxiv.2002.05651.

Hendrycks, Dan, and Thomas Dietterich. 2019. "Benchmarking Neural Network Robustness to Common Corruptions and Perturbations." *arXiv Preprint arXiv:1903.12261*.

Hendrycks, Dan, and Kevin Gimpel. 2017. "A Baseline for Detecting Misclassified and Out-of-Distribution Examples in Neural Networks." *International Conference on Learning Representations (ICLR)*.

Hennessy, J. L., and D. A. Patterson. 2011. *Computer Architecture: A Quantitative Approach*. Morgan Kaufmann.

Hennessy, John L., and David A. Patterson. 2017. *Computer Architecture: A Quantitative Approach*. 6th ed. Morgan Kaufmann.

Hennessy, John L., and David A. Patterson. 2019. "A New Golden Age for Computer Architecture." *Communications of the ACM* 62 (2): 48–60. https://doi.org/10.1145/3282307.

Hermann, Jeremy, and Mike Del Balso. 2017. *Meet Michelangelo: Uber's Machine Learning Platform*. Uber Engineering Blog.

Hestness, Joel, Sharan Narang, Newsha Ardalani, Gregory Diamos, Heewoo Jun, Hassan Kianinejad, Md. Mostofa Ali Patwary, Yang Yang, and Yanqi Zhou. 2017. "Deep Learning Scaling Is Predictable, Empirically." *arXiv Preprint arXiv:1712.00409*.

Himmelstein, Gracie, David Bates, and Li Zhou. 2022. "Examination of Stigmatizing Language in the Electronic Health Record." *JAMA Network Open* 5 (1): e2144967. https://doi.org/10.1001/jamanetworkopen.2021.44967.

Hindman, Benjamin, Andy Konwinski, Matei Zaharia, Ali Ghodsi, Anthony D. Joseph, Randy Katz, Scott Shenker, and Ion Stoica. 2011. "Mesos: A Platform for Fine-Grained Resource Sharing in the Data Center." *8th USENIX Symposium on Networked Systems Design and Implementation (NSDI 11)*, 295–308.

Hinton, G. E., N. Srivastava, A. Krizhevsky, I. Sutskever, and R. R. Salakhutdinov. 2012. "Improving Neural Networks by Preventing Co-Adaptation of Feature Detectors." *arXiv Preprint arXiv:1207.0580*.

Hinton, Geoffrey, Oriol Vinyals, and Jeff Dean. 2015. "Distilling the Knowledge in a Neural Network." *arXiv Preprint*.

Ho, Qirong, James Cipar, Henggang Cui, Seunghak Lee, Jin Kyu Kim, Phillip B. Gibbons, Garth A. Gibson, Greg Ganger, and Eric P. Xing. 2013. "More Effective Distributed ML via a Stale Synchronous Parallel Parameter Server." *Advances in Neural Information Processing Systems (NeurIPS)* 26.

Hochschild, Peter H., Paul Turner, Jeffrey C. Mogul, Rama Govindaraju, Parthasarathy Ranganathan, David E. Culler, and Amin Vahdat. 2021. "Cores That Don't Count." *Proceedings of the Workshop on Hot Topics in Operating Systems*, 9–16. https://doi.org/10.1145/3458336.3465297.

Hockney, Roger W. 1994. "The Communication Challenge for MPP: Intel Paragon and Meiko CS-2." *Parallel Computing* 20 (3): 389–98. https://doi.org/10.1016/s0167-8191(06)80021-9.

Hoffmann, Jordan, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, et al. 2022. "Training Compute-Optimal Large Language Models." *Advances in Neural Information Processing Systems (NeurIPS)* 35: 30016–30. https://doi.org/10.52202/068431-2176.

Horovod Developers. 2026. *Elastic Horovod*. Horovod documentation.

Horowitz, Mark. 2014. "1.1 Computing's Energy Problem (and What We Can Do about It)." *2014 IEEE International Solid-State Circuits Conference Digest of Technical Papers (ISSCC)*, 10–14. https://doi.org/10.1109/isscc.2014.6757323.

Hosseini, Hossein, Sreeram Kannan, Baosen Zhang, and Radha Poovendran. 2017. "Deceiving Google's Perspective API Built for Detecting Toxic Comments." *arXiv Preprint arXiv:1702.08138*, ahead of print. https://doi.org/10.48550/arXiv.1702.08138.

Houlsby, Neil, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin de Laroussilhe, Andrea Gesmundo, Mohammad Attariyan, and Sylvain Gelly. 2019. "Parameter-Efficient Transfer Learning for NLP." *Proceedings of the 36th International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 97: 2790–99.

Hsiao, Yu-Shun, Zishen Wan, Tianyu Jia, Radhika Ghosal, Abdulrahman Mahmoud, Arijit Raychowdhury, David Brooks, Gu-Yeon Wei, and Vijay Janapa Reddi. 2023. "MAVFI: An End-to-End Fault Analysis Framework with Anomaly Detection and Recovery for Micro Aerial Vehicles." *2023 Design, Automation & Test in Europe Conference & Exhibition (DATE)*, 1–6. https://doi.org/10.23919/date56975.2023.10137246.

Hsu, Liang-Ching, Ching-Yi Huang, Yen-Hsun Chuang, Ho-Wen Chen, Ya-Ting Chan, Heng Yi Teah, Tsan-Yao Chen, Chiung-Fen Chang, Yu-Ting Liu, and Yu-Min Tzou. 2016. "Accumulation of Heavy Metals and Trace Elements in Fluvial Sediments Received Effluents from Traditional and Semiconductor Industries." *Scientific Reports* 6 (1): 34250. https://doi.org/10.1038/srep34250.

Hu, Edward J., Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. "LoRA: Low-Rank Adaptation of Large Language Models." *arXiv Preprint arXiv:2106.09685*.

Hu, Qinghao, Zhisheng Ye, Zerui Wang, Guoteng Wang, Meng Zhang, Qiaoling Chen, Peng Sun, et al. 2024. "Characterization of Large Language Model Development in the Datacenter." *arXiv Preprint arXiv:2403.07648*.

Huang, Cheng, Huseyin Simitci, Yikang Xu, Aaron Ogus, Brad Calder, Parikshit Gopalan, Jin Li, and Sergey Yekhanin. 2012. "Erasure Coding in Windows Azure Storage." *2012 USENIX Annual Technical Conference (USENIX ATC 12)*, 15–26.

Huang, Y., Y. Cheng, A. Bapna, O. Firat, D. Chen, M. X. Chen, H. Lee, et al. 2019. "GPipe: Efficient Training of Giant Neural Networks Using Pipeline Parallelism." *Advances in Neural Information Processing Systems (NeurIPS)* 32: 103–12.

Hudgens, Michael G., and M. Elizabeth Halloran. 2008. "Toward Causal Inference with Interference." *Journal of the American Statistical Association* 103 (482): 832–42. https://doi.org/10.1198/016214508000000292.

Hugging Face. 2026. *Safetensors Documentation*.

Hutter, Michael, Jorn-Marc Schmidt, and Thomas Plos. 2009. "Contact-Based Fault Injections and Power Analysis on RFID Tags." *2009 European Conference on Circuit Theory and Design*, 409–12. https://doi.org/10.1109/ecctd.2009.5275012.

Hwang, Changho, Wei Cui, Yifan Xiong, Ziyue Yang, Ze Liu, Han Hu, Zilong Wang, et al. 2023. "Tutel: Adaptive Mixture-of-Experts at Scale." *Proceedings of Machine Learning and Systems (MLSys)* 5.

IEEE 802.1 Working Group. 2018. *IEEE Standard for Local and Metropolitan Area Networks—Secure Device Identity*. IEEE Std 802.1AR-2018.

IEEE Standards Association. 2019. *IEEE 754-2019: Standard for Floating-Point Arithmetic*. https://doi.org/10.1109/IEEESTD.2019.8766229.

Ignatov, Andrey, Radu Timofte, Andrei Kulik, Seungsoo Yang, Ke Wang, Felix Baum, Max Wu, Lirong Xu, and Luc Van Gool. 2019. "AI Benchmark: All about Deep Learning on Smartphones in 2019." *arXiv Preprint arXiv:1910.06663*.

Illinois General Assembly. 2020. *Artificial Intelligence Video Interview Act (Public Act 101-0260)*. 820 ILCS 42, effective January 1, 2020.

Inan, H., K. Upasani, J. Chi, R. Rungta, K. Iyer, Y. Mao, M. Tontchev, et al. 2023. "Llama Guard: LLM-Based Input-Output Safeguard for Human-AI Conversations." *arXiv Preprint arXiv:2312.06674*.

InfiniBand Trade Association. 2000. *InfiniBand Architecture Specification Volume 1*. InfiniBand Trade Association.

Intel. 2024. *Intel Gaudi 3 AI Accelerator White Paper*.

Intel Corporation. 2022. *Form 10-q for the Quarterly Period Ended July 2, 2022*.

Intel Habana Labs. 2023. *Intel Gaudi2 AI Accelerator HL-225B Mezzanine Card Datasheet*.

International Energy Agency. 2023. *World Energy Outlook 2023*. IEA. https://doi.org/10.1787/827374a6-en.

International Energy Agency. 2024a. *Data Centre Electricity Demand Estimates as a Share of Total Electricity Demand in Ireland and Virginia, 2023*. IEA chart.

International Energy Agency. 2024b. *Emissions Factors 2024*.

International Energy Agency. 2025. *Energy and AI*.

International Organization for Standardization. 2006a. *ISO 14040:2006 Environmental Management – Life Cycle Assessment – Principles and Framework*. ISO standard.

International Organization for Standardization. 2006b. *ISO 14044:2006 Environmental Management – Life Cycle Assessment – Requirements and Guidelines*. ISO standard.

International Organization for Standardization, and International Electrotechnical Commission. 2023a. *ISO/IEC 23894:2023 Information Technology – Artificial Intelligence – Guidance on Risk Management*. ISO/IEC standard.

International Organization for Standardization, and International Electrotechnical Commission. 2023b. *ISO/IEC 42001:2023 Information Technology – Artificial Intelligence – Management System*. ISO/IEC standard.

Irimia-Vladu, M. 2014. "'Green' Electronics: Biodegradable and Biocompatible Materials and Devices for Sustainable Future." *Chemical Society Reviews* 43 (2): 588–610. https://doi.org/10.1039/c3cs60235d.

Jacob, Benoit, Skirmantas Kligys, Bo Chen, Menglong Zhu, Matthew Tang, Andrew Howard, Hartwig Adam, and Dmitry Kalenichenko. 2018. "Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference." *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2704–13. https://doi.org/10.1109/cvpr.2018.00286.

Jacobs, Sam Ade, Masahiro Tanaka, Chengming Zhang, Minjia Guo, Shuaiwen Zhu, Shuai Song, Samyam Rajbhandari, Ammar Ahmad Awan, Connor Zheng, and Yuxiong He. 2023. "DeepSpeed Ulysses: System Optimizations for Enabling Training of Extreme Long Sequence Transformer Models." *arXiv Preprint arXiv:2309.14509*.

Jagielski, Matthew, Giorgio Severi, Niklas Pousette Harger, and Alina Oprea. 2021. "Subpopulation Data Poisoning Attacks." *Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security*, 3104–22. https://doi.org/10.1145/3460120.3485368.

Jain, R. 1991. *The Art of Computer Systems Performance Analysis: Techniques for Experimental Design, Measurement, Simulation, and Modeling*. John Wiley & Sons.

Janardhan, Santosh. 2021. *More Details about the October 4 Outage*. Meta Engineering postmortem.

Jeaugey, Sylvain. 2017. *NCCL 2.0*. GPU Technology Conference presentation.

Jeon, Myeongjae, Shivaram Venkataraman, Amar Phanishayee, Junjie Qian, Wencong Xiao, and Fan Yang. 2019. "Analysis of Large-Scale Multi-Tenant GPU Clusters for DNN Training Workloads." *2019 USENIX Annual Technical Conference (USENIX ATC 19)*, 947–60.

Jevons, William Stanley. 1865. *The Coal Question: An Inquiry Concerning the Progress of the Nation, and the Probable Exhaustion of Our Coal Mines*. Macmillan; Co.

Jha, A. R. 2014. *Rare Earth Materials: Properties and Applications*. CRC Press. https://doi.org/10.1201/b17045.

Jiang, A. Q., A. Sablayrolles, A. Roux, A. Mensch, B. Savary, C. Bamford, D. S. Chaplot, et al. 2024. "Mixtral of Experts." *arXiv Preprint arXiv:2401.04088*.

Jiang, Ziheng, Haibin Lin, Yinmin Zhong, Qi Huang, Yangrui Chen, Zhi Zhang, Yanghua Peng, et al. 2024. "MegaScale: Scaling Large Language Model Training to More Than 10,000 GPUs." *21st USENIX Symposium on Networked Systems Design and Implementation (NSDI 24)*, 745–60.

Jouppi, Norman P., Doe Hyun Yoon, Matthew Ashcraft, Mark Gottscho, Thomas B. Jablin, George Kurian, James Laudon, et al. 2021. "Ten Lessons from Three Generations Shaped Google's TPUv4i: Industrial Product." *2021 ACM/IEEE 48th Annual International Symposium on Computer Architecture (ISCA)*, 1–14. https://doi.org/10.1109/isca52012.2021.00010.

Jouppi, Norman P., Cliff Young, Nishant Patil, David Patterson, Gaurav Agrawal, Raminder Bajwa, Sarah Bates, et al. 2017. "In-Datacenter Performance Analysis of a Tensor Processing Unit." *Proceedings of the 44th Annual International Symposium on Computer Architecture*, ISCA '17, 1–12. https://doi.org/10.1145/3079856.3080246.

Jouppi, Norm, George Kurian, Sheng Li, Peter Ma, Rahul Nagarajan, Lifeng Nai, Nishant Patil, et al. 2023. "TPU v4: An Optically Reconfigurable Supercomputer for Machine Learning with Hardware Support for Embeddings." *Proceedings of the 50th Annual International Symposium on Computer Architecture*, 1–14. https://doi.org/10.1145/3579371.3589350.

Joye, Marc, and Michael Tunstall. 2012. *Fault Analysis in Cryptography*. Inf. Secur. Cryptography. Springer Berlin Heidelberg. https://doi.org/10.1007/978-3-642-29656-7.

Kairouz, P., and H. B. McMahan. 2021. "Advances and Open Problems in Federated Learning." *Foundations and Trends in Machine Learning* 14 (1-2): 1–210. https://doi.org/10.1561/2200000083.

Kanev, Svilen, Juan Pablo Darago, Kim Hazelwood, Parthasarathy Ranganathan, Tipp Moseley, Gu-Yeon Wei, and David Brooks. 2015. "Profiling a Warehouse-Scale Computer." *Proceedings of the 42nd Annual International Symposium on Computer Architecture*, 158–69. https://doi.org/10.1145/2749469.2750392.

Kang, Daniel, Ankit Mathur, Teja Veeramacheneni, Peter Bailis, and Matei Zaharia. 2020. "Jointly Optimizing Preprocessing and Inference for DNN-Based Visual Analytics." *Proceedings of the VLDB Endowment* 14 (2): 87–100. https://doi.org/10.14778/3425879.3425881.

Kaplan, Jared, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. "Scaling Laws for Neural Language Models." *ArXiv Preprint* abs/2001.08361.

Karger, D., E. Lehman, T. Leighton, R. Panigrahy, M. Levine, and D. Lewin. 1997. "Consistent Hashing and Random Trees: Distributed Caching Protocols for Relieving Hot Spots on the World Wide Web." *Proceedings of the Twenty-Ninth Annual ACM Symposium on Theory of Computing - STOC '97*, 654–63. https://doi.org/10.1145/258533.258660.

Karimireddy, Sai Praneeth, Quentin Rebjock, Sebastian U. Stich, and Martin Jaggi. 2019. "Error Feedback Fixes SignSGD and Other Gradient Compression Schemes." *Proceedings of the 36th International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 97: 3252–61.

Kasperkevic, Jana. 2015. "Google Says Sorry for Racist Auto-Tag in Photo App." *The Guardian*, July.

Kawazoe Aguilera, Marcos, Wei Chen, and Sam Toueg. 1997. "Heartbeat: A Timeout-Free Failure Detector for Quiescent Reliable Communication." In *Lect. Notes Comput. Sci.* Springer Berlin Heidelberg. https://doi.org/10.1007/bfb0030680.

Kim, B., M. Wattenberg, J. Gilmer, C. Cai, J. Wexler, F. Viegas, and R. Sayres. 2018. "Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors (TCAV)." *Proceedings of the 35th International Conference on Machine Learning (ICML)* 80: 2668–77.

Kim, J., W. J. Dally, S. Scott, and D. Abts. 2008. "Technology-Driven, Highly-Scalable Dragonfly Topology." *2008 International Symposium on Computer Architecture*, 77–88. https://doi.org/10.1109/isca.2008.19.

Kim, Sunju, Chungsik Yoon, Seunghon Ham, Jihoon Park, Ohun Kwon, Donguk Park, Sangjun Choi, Seungwon Kim, Kwonchul Ha, and Won Kim. 2018. "Chemical Use in the Semiconductor Manufacturing Industry." *International Journal of Occupational and Environmental Health* 24 (3-4): 109–18. https://doi.org/10.1080/10773525.2018.1519957.

Kingma, Diederik P., and Jimmy Ba. 2015. "Adam: A Method for Stochastic Optimization." *3rd International Conference on Learning Representations (ICLR)*.

Kirkpatrick, J., R. Pascanu, N. Rabinowitz, J. Veness, G. Desjardins, A. A. Rusu, K. Milan, et al. 2017. "Overcoming Catastrophic Forgetting in Neural Networks." *Proceedings of the National Academy of Sciences* 114 (13): 3521–26. https://doi.org/10.1073/pnas.1611835114.

Kleinberg, Jon, Sendhil Mullainathan, and Manish Raghavan. 2017. "Inherent Trade-Offs in the Fair Determination of Risk Scores." *8th Innovations in Theoretical Computer Science Conference (ITCS 2017)*, Leibniz international proceedings in informatics (LIPIcs), vol. 67: 43:1–23. https://doi.org/10.4230/LIPIcs.ITCS.2017.43.

Kleinrock, L. 1975. *Queueing Systems, Volume 1: Theory*. Wiley-Interscience.

Kleppmann, Martin. 2016. *Designing Data-Intensive Applications: The Big Ideas Behind Reliable, Scalable, and Maintainable Systems*. O'Reilly Media.

Klutke, G., P. C. Kiessler, and M. A. Wortman. 2003. "A Critical Look at the Bathtub Curve." *IEEE Transactions on Reliability* 52 (1): 125–29. https://doi.org/10.1109/tr.2002.804492.

Ko, Y. 2021. "Characterizing System-Level Masking Effects Against Soft Errors." *Electronics* 10 (18): 2286. https://doi.org/10.3390/electronics10182286.

Kocher, Paul, Jann Horn, Anders Fogh, Daniel Genkin, Daniel Gruss, Werner Haas, Mike Hamburg, et al. 2019. "Spectre Attacks: Exploiting Speculative Execution." *2019 IEEE Symposium on Security and Privacy (SP)*, 1–19. https://doi.org/10.1109/sp.2019.00002.

Kocher, Paul, Joshua Jaffe, and Benjamin Jun. 1999. "Differential Power Analysis." In *Advances in Cryptology — CRYPTO' 99*. Springer Berlin Heidelberg. https://doi.org/10.1007/3-540-48405-1_25.

Kocher, Paul, Joshua Jaffe, Benjamin Jun, and Pankaj Rohatgi. 2011. "Introduction to Differential Power Analysis." *Journal of Cryptographic Engineering* 1 (1): 5–27. https://doi.org/10.1007/s13389-011-0006-y.

Koh, Pang Wei, and Percy Liang. 2017. "Understanding Black-Box Predictions via Influence Functions." *International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 70: 1885–94.

Koh, Pang Wei, Thao Nguyen, Yew Siang Tang, Stephen Mussmann, Emma Pierson, Been Kim, and Percy Liang. 2020. "Concept Bottleneck Models." *Proceedings of the 37th International Conference on Machine Learning, ICML 2020, 13-18 July 2020, Virtual Event*, Proceedings of machine learning research, vol. 119: 5338–48.

Kohavi, Ron, Roger Longbotham, Dan Sommerfield, and Randal M. Henne. 2009. "Controlled Experiments on the Web: Survey and Practical Guide." *Data Mining and Knowledge Discovery* 18 (1): 140–81. https://doi.org/10.1007/s10618-008-0114-1.

Kohavi, Ron, Diane Tang, and Ya Xu. 2020. *Trustworthy Online Controlled Experiments: A Practical Guide to A/B Testing*. Cambridge University Press. https://doi.org/10.1017/9781108653985.

Kokolis, Apostolos, Michael Kuchnik, John Hoffman, Adithya Kumar, Parth Malani, Faye Ma, Zach DeVito, Shubho Sengupta, Kalyan Saladi, and Carole-Jean Wu. 2025. "Revisiting Reliability in Large-Scale Machine Learning Research Clusters." *2025 IEEE International Symposium on High Performance Computer Architecture (HPCA)*, 1259–74. https://doi.org/10.1109/hpca61900.2025.00096.

Konečný, J., H. B. McMahan, D. Ramage, and P. Richtárik. 2016. "Federated Optimization: Distributed Machine Learning for on-Device Intelligence." *CoRR* abs/1610.02527.

Koren, Yehuda, Robert Bell, and Chris Volinsky. 2009. "Matrix Factorization Techniques for Recommender Systems." *Computer* 42 (8): 30–37. https://doi.org/10.1109/mc.2009.263.

Korosec, Kirsten. 2019. *Tesla Sues Former Employees, Zoox for Alleged Trade Secret Theft*. TechCrunch.

Korthikanti, Vijay Anand, Jared Casper, Sangkug Lym, Lawrence McAfee, Michael Andersch, Mohammad Shoeybi, and Bryan Catanzaro. 2023. "Reducing Activation Recomputation in Large Transformer Models." *Proceedings of Machine Learning and Systems (MLSys)* 5: 341–53.

Krishna, Kalpesh, Gaurav Singh Tomar, Ankur P. Parikh, Nicolas Papernot, and Mohit Iyyer. 2020. "Thieves on Sesame Street! Model Extraction of BERT-Based APIs." *International Conference on Learning Representations (ICLR)*.

Krishnamoorthi, Raghuraman. 2018. "Quantizing Deep Convolutional Networks for Efficient Inference: A Whitepaper." *arXiv Preprint arXiv:1806.08342* abs/1806.08342.

Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. 2012. "ImageNet Classification with Deep Convolutional Neural Networks." *Advances in Neural Information Processing Systems (NeurIPS)* 25.

Kung, H. T. 1982. "Why Systolic Architectures?" *Computer* 15 (1): 37–46. https://doi.org/10.1109/mc.1982.1653825.

Kung, Hsiang Tsung, and Charles E Leiserson. 1979. "Systolic Arrays (for VLSI)." *Sparse Matrix Proceedings 1978* 1: 256–82.

Kurakin, Alexey, Ian Goodfellow, and Samy Bengio. 2017. "Adversarial Examples in the Physical World." *International Conference on Learning Representations (ICLR), Workshop Track*, 99–112. https://doi.org/10.1201/9781351251389-8.

Kusano, Kristofer D., John M. Scanlon, Yin-Hsiu Chen, Timothy L. McMurry, Ruoshu Chen, Tilia Gode, and Trent Victor. 2024. "Comparison of Waymo Rider-Only Crash Data to Human Benchmarks at 7.1 Million Miles." *Traffic Injury Prevention* 25 (sup1). https://doi.org/10.1080/15389588.2024.2380786.

Kusner, Matt J., Joshua Loftus, Chris Russell, and Ricardo Silva. 2017. "Counterfactual Fairness." *Advances in Neural Information Processing Systems* 30: 4066–76.

Kwon, Woosuk, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. 2023. "Efficient Memory Management for Large Language Model Serving with PagedAttention." *Proceedings of the 29th Symposium on Operating Systems Principles*, 611–26. https://doi.org/10.1145/3600006.3613165.

Kwon, Y. D., R. Li, S. I. Venieris, J. Chauhan, and N. D. Lane. 2024. "TinyTrain: Resource-Aware Task-Adaptive Sparse Training of DNNs at the Data-Scarce Edge." *Proceedings of the 41st International Conference on Machine Learning (ICML)*.

Kwon, Youngeun, and Minsoo Rhu. 2018. "TensorDIMM: A Practical Near-Memory Processing Architecture for Embeddings and Tensor Operations in Deep Learning." *Proceedings of the 52nd Annual IEEE/ACM International Symposium on Microarchitecture*, 740–53. https://doi.org/10.1145/3352460.3358284.

Lacoste, Alexandre, Alexandra Luccioni, Victor Schmidt, and Thomas Dandres. 2019. "Quantifying the Carbon Emissions of Machine Learning." *arXiv Preprint arXiv:1910.09700*.

Lai, Liangzhen, Naveen Suda, and Vikas Chandra. 2018. "CMSIS-NN: Efficient Neural Network Kernels for Arm Cortex-M CPUs." *ArXiv Preprint* abs/1801.06601.

Lakshminarayanan, Balaji, Alexander Pritzel, and Charles Blundell. 2017. "Simple and Scalable Predictive Uncertainty Estimation Using Deep Ensembles." *Advances in Neural Information Processing Systems (NeurIPS)* 30.

Lamport, Leslie, Robert Shostak, and Marshall Pease. 1982. "The Byzantine Generals Problem." *ACM Transactions on Programming Languages and Systems* 4 (3): 382–401. https://doi.org/10.1145/357172.357176.

Langston, Jennifer. 2020. *Microsoft Announces New Supercomputer, Lays Out Vision for Future AI Work*. Microsoft Source.

Lannelongue, Löic, Jason Grealey, and Michael Inouye. 2021. "Green Algorithms: Quantifying the Carbon Footprint of Computation." *Advanced Science* 8 (12): 2100707. https://doi.org/10.1002/advs.202100707.

Lazer, David, Ryan Kennedy, Gary King, and Alessandro Vespignani. 2014. "The Parable of Google Flu: Traps in Big Data Analysis." *Science* 343 (6176): 1203–5. https://doi.org/10.1126/science.1248506.

Lee, Peter. 2016. *Learning from Tay's Introduction*. The Official Microsoft Blog.

Lei, Jing, Max G'Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. 2018. "Distribution-Free Predictive Inference for Regression." *Journal of the American Statistical Association* 113 (523): 1094–111. https://doi.org/10.1080/01621459.2017.1307116.

Leiserson, Charles E. 1985. "Fat-Trees: Universal Networks for Hardware-Efficient Supercomputing." *IEEE Transactions on Computers* C-34 (10): 892–901. https://doi.org/10.1109/tc.1985.6312192.

Lepikhin, Dmitry, HyoukJoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. 2021. "GShard: Scaling Giant Models with Conditional Computation and Automatic Sharding." *International Conference on Learning Representations (ICLR)*.

Leviathan, Yaniv, Matan Kalman, and Yossi Matias. 2023. "Fast Inference from Transformers via Speculative Decoding." *International Conference on Machine Learning (ICML)*, 19274–86.

Lewis, Patrick, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, et al. 2020. "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks." *Advances in Neural Information Processing Systems* 33: 9459–74.

Li, Chengwei. 2020. *Estimating the Training Cost of GPT-3*.

Li, Jie, George Michelogiannakis, Brandon Cook, Dulanya Cooray, and Yong Chen. 2023. "Analyzing Resource Utilization in an HPC System: A Case Study of NERSC's Perlmutter." *Lect. Notes Comput. Sci.*, Lecture notes in computer science, vol. 13948: 297–316. https://doi.org/10.1007/978-3-031-32041-5_16.

Li, M., D. G. Andersen, J. W. Park, A. J. Smola, A. Ahmed, V. Josifovski, J. Long, E. J. Shekita, and B.-Y. Su. 2014. "Scaling Distributed Machine Learning with the Parameter Server." *Proceedings of the 2014 International Conference on Big Data Science and Computing*, 583–98. https://doi.org/10.1145/2640087.2644155.

Li, Shen, Yanli Zhao, Rohan Varma, Omkar Salpekar, Pieter Noordhuis, Teng Li, Adam Paszke, et al. 2020. "PyTorch Distributed: Experiences on Accelerating Data Parallel Training." *Proceedings of the VLDB Endowment* 13 (12): 3005–18. https://doi.org/10.14778/3415478.3415530.

Li, Tian, Anit Kumar Sahu, Ameet Talwalkar, and Virginia Smith. 2020. "Federated Learning: Challenges, Methods, and Future Directions." *IEEE Signal Processing Magazine* 37 (3): 50–60. https://doi.org/10.1109/msp.2020.2975749.

Li, Yuliang, Rui Miao, Hongqiang Harry Liu, Yan Zhuang, Fei Feng, Lingbo Tang, Zheng Cao, et al. 2019. "HPCC: High Precision Congestion Control." *Proceedings of the ACM Special Interest Group on Data Communication (SIGCOMM)*, 44–58. https://doi.org/10.1145/3341302.3342085.

Lifka, David A. 1995. "The ANL/IBM SP Scheduling System." In *Job Scheduling Strategies for Parallel Processing*, edited by Dror G. Feitelson and Larry Rudolph, vol. 949. Lecture Notes in Computer Science. Springer. https://doi.org/10.1007/3-540-60153-8_35.

Lin, Ji, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. 2024. "AWQ: Activation-Aware Weight Quantization for on-Device LLM Compression and Acceleration." *Proceedings of Machine Learning and Systems* 6: 87–100.

Lin, Yujun, Song Han, Huizi Mao, Yu Wang, and William J Dally. 2018. "Deep Gradient Compression: Reducing the Communication Bandwidth for Distributed Training." *International Conference on Learning Representations (ICLR)*.

Linux man-pages project. 2026. *Perfquery(8): Query InfiniBand Port Counters on a Single Port*.

linux-rdma project. 2026. *perftest: InfiniBand Verbs Performance Tests*.

Lipp, M., M. Schwarz, D. Gruss, T. Prescher, W. Haas, J. Horn, S. Mangard, et al. 2018. "Meltdown: Reading Kernel Memory from User Space." *Proceedings of the 27th USENIX Security Symposium (USENIX Security)*, 973–90.

Lipton, Zachary, Yu-Xiang Wang, and Alexander Smola. 2018. "Detecting and Correcting for Label Shift with Black Box Predictors." *International Conference on Machine Learning (ICML)*, 3122–30.

Little, John D. C. 1961. "A Proof for the Queuing Formula: $L = \lambda W$." *Operations Research* 9 (3): 383–87. https://doi.org/10.1287/opre.9.3.383.

Liu, Hao, Matei Zaharia, and Pieter Abbeel. 2023. "Ring Attention with Blockwise Transformers for Near-Infinite Context." *arXiv Preprint arXiv:2310.01889*, ahead of print. https://doi.org/10.48550/arXiv.2310.01889.

Liu, W., Y. Xi, J. Qin, F. Sun, B. Chen, W. Zhang, R. Zhang, and R. Tang. 2022. "Neural Re-Ranking in Multi-Stage Recommender Systems: A Review." *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence*, 5512–20. https://doi.org/10.24963/ijcai.2022/771.

Luccioni, Alexandra Sasha, Sylvain Viguier, and Anne-Laure Ligozat. 2023. "Estimating the Carbon Footprint of BLOOM, a 176B Parameter Language Model." *Journal of Machine Learning Research* 24 (253): 1–15.

Luccioni, Sasha, Boris Gamazaychikov, Theo Alves da Costa, and Emma Strubell. 2025. "Misinformation by Omission: The Need for More Environmental Transparency in AI." *arXiv Preprint arXiv:2506.15572*.

Lundberg, Scott M., and Su-In Lee. 2017. "A Unified Approach to Interpreting Model Predictions." *Advances in Neural Information Processing Systems (NeurIPS)* 30: 4765–74.

Lustre Wiki. 2017. *lfs setstripe: Set Striping Pattern of a File*. Lustre Wiki manual page.

Lustre Wiki. 2019. *Configuring Lustre File Striping*. Lustre Wiki.

Ma, Jeffrey, Hengzhi Pei, Leonard Lausen, and George Karypis. 2025. *Understanding Silent Data Corruption in LLM Training*. https://doi.org/10.48550/arXiv.2502.12340.

Ma, Jeffrey, Alan Tu, Yiling Chen, and Vijay Janapa Reddi. 2024. "FedStaleWeight: Buffered Asynchronous Federated Learning with Fair Aggregation via Staleness Reweighting." *arXiv Preprint arXiv:2406.02877*.

Ma, Xiaoyu, and David Patterson. 2026. *Challenges and Research Directions for Large Language Model Inference Hardware*. No. 5. Vol. 59. https://doi.org/10.1109/mc.2026.3652916.

Madry, Aleksander, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. 2018. "Towards Deep Learning Models Resistant to Adversarial Attacks." *International Conference on Learning Representations (ICLR)*.

Mahajan, Kshiteej, Arjun Balasubramanian, Arjun Singhvi, Shivaram Venkataraman, Aditya Akella, Amar Phanishayee, and Shuchi Chawla. 2020. "Themis: Fair and Efficient GPU Cluster Scheduling." *17th USENIX Symposium on Networked Systems Design and Implementation (NSDI 20)*, 289–304.

Mahmoud, A., N. Aggarwal, A. Nobbe, J. R. S. Vicarte, S. V. Adve, C. W. Fletcher, I. Frosio, and S. K. S. Hari. 2020. "PyTorchFI: A Runtime Perturbation Tool for DNNs." *2020 50th Annual IEEE/IFIP International Conference on Dependable Systems and Networks Workshops (DSN-w)*, 25–31. https://doi.org/10.1109/dsn-w50199.2020.00014.

Malkov, Yury A., and Dmitry A. Yashunin. 2020. "Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs." *IEEE Transactions on Pattern Analysis and Machine Intelligence* 42 (4): 824–36. https://doi.org/10.1109/TPAMI.2018.2889473.

Masanet, Eric, Arman Shehabi, Nuoa Lei, Sarah Smith, and Jonathan Koomey. 2020. "Recalibrating Global Data Center Energy-Use Estimates." *Science* 367 (6481): 984–86. https://doi.org/10.1126/science.aba3758.

Mattson, Peter, Christine Cheng, Cody Coleman, Greg Diamos, Paulius Micikevicius, David Patterson, Hanlin Tang, et al. 2020. "MLPerf Training Benchmark." *Proceedings of Machine Learning and Systems (MLSys)* 2: 336–49.

McCandlish, Sam, Jared Kaplan, Dario Amodei, and OpenAI Dota Team. 2018. "An Empirical Model of Large-Batch Training." *arXiv Preprint arXiv:1812.06162*, ahead of print. https://doi.org/10.48550/arXiv.1812.06162.

McCarthy, J. 1981. "Epistemological Problems of Artificial Intelligence." In *Readings in Artificial Intelligence*. Elsevier. https://doi.org/10.1016/b978-0-934613-03-3.50035-0.

McMahan, Brendan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera y Arcas. 2017. "Communication-Efficient Learning of Deep Networks from Decentralized Data." *International Conference on Artificial Intelligence and Statistics (AISTATS)*, Proceedings of machine learning research, vol. 54: 1273–82.

Message Passing Interface Forum. 2015. *MPI: A Message-Passing Interface Standard, Version 3.1*.

Meta AI. 2024. *Llama 3.1 Model Card*.

Meta AI Research. 2022. *OPT-175B Training Logbook*. Companion repository to the OPT-175B model release.

Meta Engineering. 2024. *Building Meta's GenAI Infrastructure*. Engineering at Meta Blog.

Micikevicius, Paulius, Dusan Stosic, Neil Burgess, Marius Cornea, Pradeep Dubey, Richard Grisenthwaite, Sangwon Ha, et al. 2022. "FP8 Formats for Deep Learning." *arXiv Preprint arXiv:2209.05433*.

Microsoft. 2025. *Inside the World's Most Powerful AI Datacenter*. The Official Microsoft Blog.

Miller, C. 2019. "Lessons Learned from Hacking a Car." *IEEE Design & Test* 36 (6): 7–9. https://doi.org/10.1109/mdat.2018.2863106.

Miller, Charlie, and Chris Valasek. 2015. *Remote Exploitation of an Unaltered Passenger Vehicle*. Black Hat USA 2015 Whitepaper.

Mironov, I. 2017. "Rényi Differential Privacy." *2017 IEEE 30th Computer Security Foundations Symposium (CSF)*, 263–75. https://doi.org/10.1109/csf.2017.11.

Mitchell, Margaret, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. 2019. "Model Cards for Model Reporting." *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 220–29. https://doi.org/10.1145/3287560.3287596.

Mittal, Radhika, Alexander Shpiner, Aurojit Panda, Eitan Zahavi, Arvind Krishnamurthy, Sylvia Ratnasamy, and Scott Shenker. 2018. "Revisiting Network Support for RDMA." *Proceedings of the 2018 Conference of the ACM Special Interest Group on Data Communication*, 313–26. https://doi.org/10.1145/3230543.3230557.

Mitzenmacher, M. 2001. "The Power of Two Choices in Randomized Load Balancing." *IEEE Transactions on Parallel and Distributed Systems* 12 (10): 1094–104. https://doi.org/10.1109/71.963420.

MLCommons. 2024. *MLPerf Storage V1.0 Results*.

MLCommons. 2025a. *MLPerf Inference: Datacenter Benchmark Results, Rounds V3.1 to V5.0*.

MLCommons. 2025b. *MLPerf Training Benchmark Results, Rounds V3.1 to V5.1*.

ML.ENERGY Initiative, Jae-Won Chung, Jiachen Gu, Zhewei Yan, Tony Lo, Yejin Park, and Mosharaf Chowdhury. 2023. *The ML.ENERGY Leaderboard*. SymbioticLab, University of Michigan.

Mnih, Volodymyr, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, et al. 2015. "Human-Level Control Through Deep Reinforcement Learning." *Nature* 518 (7540): 529–33. https://doi.org/10.1038/nature14236.

Mohanram, K., and N. A. Touba. 2003. "Partial Error Masking to Reduce Soft Error Failure Rate in Logic Circuits." *Proceedings. 18th IEEE International Symposium on Defect and Fault Tolerance in VLSI Systems*, 433–40. https://doi.org/10.1109/dftvs.2003.1250141.

Monyei, C. G., and Kirsten E. H. Jenkins. 2018. "Electrons Have No Identity: Setting Right Misrepresentations in Google and Apple's Clean Energy Purchasing." *Energy Research & Social Science* 46: 48–51. https://doi.org/10.1016/j.erss.2018.06.015.

Moosavi-Dezfooli, Seyed-Mohsen, Alhussein Fawzi, and Pascal Frossard. 2015. "DeepFool: A Simple and Accurate Method to Fool Deep Neural Networks." *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* 2016: 2574–82. https://doi.org/10.1109/cvpr.2016.282.

Mu, Norman, and Justin Gilmer. 2019. *MNIST-C: A Robustness Benchmark for Computer Vision*.

Mukherjee, S. S., J. Emer, and S. K. Reinhardt. 2005. "The Soft Error Problem: An Architectural Perspective." *11th International Symposium on High-Performance Computer Architecture*, 243–47. https://doi.org/10.1109/hpca.2005.37.

Nabi, Razieh, and Ilya Shpitser. 2018. "Fair Inference on Outcomes." *Proceedings of the AAAI Conference on Artificial Intelligence* 32. https://doi.org/10.1609/aaai.v32i1.11553.

Narayanan, Arvind, and Vitaly Shmatikov. 2006. "How to Break Anonymity of the Netflix Prize Dataset." In *CoRR*, abs/cs/0610105.

Narayanan, Arvind, and Vitaly Shmatikov. 2008. "Robust de-Anonymization of Large Sparse Datasets." *2008 IEEE Symposium on Security and Privacy (SP 2008)*, 111–25. https://doi.org/10.1109/sp.2008.33.

Narayanan, Deepak, Aaron Harlap, Amar Phanishayee, Vivek Seshadri, Nikhil R. Devanur, Gregory R. Ganger, Phillip B. Gibbons, and Matei Zaharia. 2019. "PipeDream: Generalized Pipeline Parallelism for DNN Training." *Proceedings of the 27th ACM Symposium on Operating Systems Principles*, 1–15. https://doi.org/10.1145/3341301.3359646.

Narayanan, Deepak, Mohammad Shoeybi, Jared Casper, Patrick LeGresley, Mostofa Patwary, Vijay Korthikanti, Dmitri Vainbrand, et al. 2021. "Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM." *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, 1–15. https://doi.org/10.1145/3458817.3476209.

NASA Mars Program Independent Assessment Team. 2000. *Report on the Loss of the Mars Polar Lander and Deep Space 2 Missions*. National Aeronautics; Space Administration.

Nasr, Milad, Reza Shokri, and Amir Houmansadr. 2019. "Comprehensive Privacy Analysis of Deep Learning: Passive and Active White-Box Inference Attacks Against Centralized and Federated Learning." *2019 IEEE Symposium on Security and Privacy (SP)*, 739–53. https://doi.org/10.1109/sp.2019.00065.

National Transportation Safety Board. 2017. *Collision Between a Car Operating with Automated Vehicle Control Systems and a Tractor-Semitrailer Truck Near Williston, Florida, May 7, 2016*. HAR-17/02. National Transportation Safety Board.

National Transportation Safety Board. 2019. *Collision Between Vehicle Controlled by Developmental Automated Driving System and Pedestrian, Tempe, Arizona, March 18, 2018*. HAR-19/03. National Transportation Safety Board.

Naumov, Maxim, Dheevatsa Mudigere, Hao-Jun Michael Shi, Jianyu Huang, Narayanan Sundaraman, Jongsoo Park, Xiaodong Wang, et al. 2019. "Deep Learning Recommendation Model for Personalization and Recommendation Systems." *arXiv Preprint arXiv:1906.00091*.

Netflix Technology Blog. 2017. *Innovating Faster on Personalization Algorithms at Netflix Using Interleaving*. Netflix Tech Blog.

Neumann, John von. 1928. "Zur Theorie Der Gesellschaftsspiele." *Mathematische Annalen* 100 (1): 295–320. https://doi.org/10.1007/bf01448847.

New York State Department of Financial Services. 2021. *Report on Apple Card Investigation*. New York State Department of Financial Services.

Ngo, Richard, Lawrence Chan, and Sören Mindermann. 2022. "The Alignment Problem from a Deep Learning Perspective." *ArXiv Preprint* abs/2209.00626.

Nguyen, John, Kshitiz Malik, Hongyuan Zhan, Ashkan Yousefpour, Michael Rabbat, Mani Malek, and Dzmitry Huba. 2021. "Federated Learning with Buffered Asynchronous Aggregation." *Proceedings of Machine Learning Research (AISTATS)* 164.

NVIDIA. 2013. *Tesla K20X GPU Accelerator Board Specification*.

NVIDIA. 2023. *NVIDIA DGX SuperPOD: Next Generation Scalable Infrastructure for AI Leadership*. RA-11333-001 V11. NVIDIA.

NVIDIA. 2024. *PCF Summary for NVIDIA HGX B200*.

NVIDIA. 2025. *Inside NVIDIA Blackwell Ultra: The Chip Powering the AI Factory Era*.

NVIDIA. 2026a. *Exemplar Performance: NCCL Microbenchmarks*.

NVIDIA. 2026b. *GPUDirect RDMA*.

NVIDIA. 2026c. *NCCL User Guide*.

NVIDIA. 2026d. *NVIDIA CUDA Multi-Process Service*.

NVIDIA. 2026e. *NVIDIA Multi-Instance GPU User Guide*.

NVIDIA. 2026f. *NVIDIA Quantum-X800 InfiniBand Platform*. NVIDIA product documentation.

NVIDIA Corporation. 2017. *NVIDIA Tesla V100 GPU Architecture*. NVIDIA Whitepaper.

NVIDIA Corporation. 2020. *NVIDIA A100 Tensor Core GPU Architecture*. NVIDIA Whitepaper, V1.0.

NVIDIA Corporation. 2025. *Product Carbon Footprint Summary for NVIDIA HGX H100*. NVIDIA Datasheet.

Nygard, Michael T. 2007. *Release It!: Design and Deploy Production-Ready Software*. Pragmatic Bookshelf.

O'Connor, Mike, Niladrish Chatterjee, Donghyuk Lee, John Wilson, Aditya Agrawal, Stephen W. Keckler, and William J. Dally. 2017. "Fine-Grained DRAM: Energy-Efficient DRAM for Extreme Bandwidth Systems." *Proceedings of the 50th Annual IEEE/ACM International Symposium on Microarchitecture*, 41–54. https://doi.org/10.1145/3123939.3124545.

Obermeyer, Ziad, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. 2019. "Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations." *Science* 366 (6464): 447–53. https://doi.org/10.1126/science.aax2342.

Ofqual. 2020. *Awarding GCSE, AS, A Level, Advanced Extension Awards and Extended Project Qualifications in Summer 2020: Interim Report*. Office of Qualifications; Examinations Regulation.

Olah, C., N. Cammarata, L. Schubert, G. Goh, M. Petrov, and S. Carter. 2020. "Zoom in: An Introduction to Circuits." *Distill* 5 (3): e00024–001. https://doi.org/10.23915/distill.00024.001.

Oliynyk, Daryna, Rudolf Mayer, and Andreas Rauber. 2023. "I Know What You Trained Last Summer: A Survey on Stealing Machine Learning Models and Defences." *ACM Computing Surveys* 55 (14s): 1–41. https://doi.org/10.1145/3595292.

Olston, Christopher, Noah Fiedel, Kiril Gorovoy, Jeremiah Harmsen, Li Lao, Fangwei Li, Vinu Rajashekhar, Sukriti Ramesh, and Jordan Soyke. 2017. "TensorFlow-Serving: Flexible, High-Performance ML Serving." *CoRR* abs/1712.06139. https://doi.org/10.48550/arXiv.1712.06139.

OpenAI. 2023. *March 20 ChatGPT Outage: Here's What Happened*. OpenAI incident report.

OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, et al. 2023. "GPT-4 Technical Report." *arXiv Preprint arXiv:2303.08774*, ahead of print. https://doi.org/10.48550/arXiv.2303.08774.

Oprea, Alina, Anoop Singhal, and Apostol Vassilev. 2022. "Poisoning Attacks Against Machine Learning: Can Machine Learning Be Trustworthy?" *Computer* 55 (11): 94–99. https://doi.org/10.1109/mc.2022.3190787.

Orekondy, Tribhuvanesh, Bernt Schiele, and Mario Fritz. 2019. "Knockoff Nets: Stealing Functionality of Black-Box Models." *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 4949–58. https://doi.org/10.1109/cvpr.2019.00509.

Organisation for Economic Co-operation and Development. 2024. *Recommendation of the Council on Artificial Intelligence*. OECD Legal Instruments.

Ousterhout, John K. 1982. "Scheduling Techniques for Concurrent Systems." *Proceedings of the 3rd International Conference on Distributed Computing Systems (ICDCS)*, 22–30.

Ouyang, L., J. Wu, X. Jiang, D. Almeida, C. L. Wainwright, P. Mishkin, C. Zhang, et al. 2022. "Training Language Models to Follow Instructions with Human Feedback." *Advances in Neural Information Processing Systems (NeurIPS)* 35: 27730–44. https://doi.org/10.52202/068431-2011.

OVHcloud. 2021. *Strasbourg Datacentre: Latest Information*. OVHcloud incident update.

Pan, Sinno Jialin, and Qiang Yang. 2010. "A Survey on Transfer Learning." *IEEE Transactions on Knowledge and Data Engineering* 22 (10): 1345–59. https://doi.org/10.1109/tkde.2009.191.

Papernot, Nicolas, Patrick McDaniel, and Ian Goodfellow. 2016. "Transferability in Machine Learning: From Phenomena to Black-Box Attacks Using Adversarial Samples." *arXiv Preprint arXiv:1605.07277*.

Papernot, Nicolas, Patrick McDaniel, Xi Wu, Somesh Jha, and Ananthram Swami. 2016. "Distillation as a Defense to Adversarial Perturbations Against Deep Neural Networks." *2016 IEEE Symposium on Security and Privacy (SP)*, 582–97. https://doi.org/10.1109/sp.2016.41.

Papernot, N., P. McDaniel, S. Jha, M. Fredrikson, Z. Berkay Celik, and A. Swami. 2016. "The Limitations of Deep Learning in Adversarial Settings." *2016 IEEE European Symposium on Security and Privacy (EuroS&P)*, 372–87. https://doi.org/10.1109/eurosp.2016.36.

Parrish, Alicia, Hannah Rose Kirk, Jessica Quaye, Charvi Rastogi, Max Bartolo, Oana Inel, Juan Ciro, et al. 2023. "Adversarial Nibbler: A Data-Centric Challenge for Improving the Safety of Text-to-Image Models." *ArXiv Preprint* abs/2305.14384.

Patarasuk, Pitch, and Xin Yuan. 2009. "Bandwidth Optimal All-Reduce Algorithms for Clusters of Workstations." *Journal of Parallel and Distributed Computing* 69 (2): 117–24. https://doi.org/10.1016/j.jpdc.2008.09.002.

Patel, Dylan, and Gerald Wong. 2023. *GPT-4 Architecture, Infrastructure, Training Dataset, Costs, Vision, MoE*. SemiAnalysis Blog.

Patel, Pratyush, Esha Choukse, Chaojie Zhang, Aashaka Shah, Íñigo Goiri, Saeed Maleki, and Ricardo Bianchini. 2024. "Splitwise: Efficient Generative LLM Inference Using Phase Splitting." *2024 ACM/IEEE 51st Annual International Symposium on Computer Architecture (ISCA)*, 118–32. https://doi.org/10.1109/isca59077.2024.00019.

Pati, Suchita, Shaizeen Aga, Mahzabeen Islam, Nuwan Jayasena, and Matthew D. Sinclair. 2024. "T3: Transparent Tracking & Triggering for Fine-Grained Overlap of Compute & Collectives." *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2*, 1146–64. https://doi.org/10.1145/3620665.3640410.

Patterson, David A., Garth Gibson, and Randy H. Katz. 1988. "A Case for Redundant Arrays of Inexpensive Disks (RAID)." *Proceedings of the 1988 ACM SIGMOD International Conference on Management of Data*, 109–16. https://doi.org/10.1145/50202.50214.

Patterson, David, Joseph Gonzalez, Urs Holzle, Quoc Le, Chen Liang, Lluis-Miquel Munguia, Daniel Rothchild, David R. So, Maud Texier, and Jeff Dean. 2022. "The Carbon Footprint of Machine Learning Training Will Plateau, Then Shrink." *Computer* 55 (7): 18–28. https://doi.org/10.1109/mc.2022.3148714.

Patterson, David, Joseph Gonzalez, Quoc Le, Chen Liang, Lluis-Miquel Munguia, Daniel Rothchild, David So, Maud Texier, and Jeff Dean. 2021. "Carbon Emissions and Large Neural Network Training." *arXiv Preprint arXiv:2104.10350*.

Perez, Fábio, and Ian Ribeiro. 2022. "Ignore Previous Prompt: Attack Techniques for Language Models." *arXiv Preprint arXiv:2211.09527*.

Personal Information Protection Commission, Japan. 2023. *Act on the Protection of Personal Information*. Personal Information Protection Commission, Japan.

Peters, D., R. A. Calvo, and R. M. Ryan. 2018. "Designing for Motivation, Engagement and Wellbeing in Digital Experience." *Frontiers in Psychology* 9: 797. https://doi.org/10.3389/fpsyg.2018.00797.

Peterson, W. W., and D. T. Brown. 1961. "Cyclic Codes for Error Detection." *Proceedings of the IRE* 49 (1): 228–35. https://doi.org/10.1109/jrproc.1961.287814.

Pilz, Konstantin F., James Sanders, Robi Rahman, and Lennart Heim. 2025. "Trends in AI Supercomputers." *arXiv Preprint arXiv:2504.16026*.

Plank, James S. 1997. "A Tutorial on Reed-Solomon Coding for Fault-Tolerance in RAID-Like Systems." *Software: Practice and Experience* 27 (9): 995–1012. https://doi.org/10.1002/(SICI)1097-024X(199709)27:9<995::AID-SPE111>3.0.CO;2-6.

Polyzotis, Neoklis, Sudip Roy, Steven Euijong Whang, and Martin Zinkevich. 2017. "Data Management Challenges in Production Machine Learning." *Proceedings of the 2017 ACM International Conference on Management of Data*, 1723–26. https://doi.org/10.1145/3035918.3054782.

Pont, Michael J, and Royan HL Ong. 2002. "Using Watchdog Timers to Improve the Reliability of Single-Processor Embedded Systems: Seven New Patterns and a Case Study." *Proceedings of the First Nordic Conference on Pattern Languages of Programs*, 159–200.

Pope, Reiner, Sholto Douglas, Aakanksha Chowdhery, Jacob Devlin, James Bradbury, Jonathan Heek, Kefan Xiao, Shivani Agrawal, and Jeff Dean. 2023. "Efficiently Scaling Transformer Inference." *Proceedings of Machine Learning and Systems (MLSys)* 5: 606–24.

Prakash, Shvetank, Matthew Stewart, Colby Banbury, Mark Mazumder, Pete Warden, Brian Plancher, and Vijay Janapa Reddi. 2023. "Is TinyML Sustainable? Assessing the Environmental Impacts of Machine Learning on Microcontrollers." *ArXiv Preprint* abs/2301.11899.

Present, Elaina, Pieter Gagnon, Eric J. H. Wilson, Noel Merket, Philip R. White, and Scott Horowitz. 2024. *Choosing the Best Carbon Factor for the Job: Exploring Available Carbon Emissions Factors and the Impact of Factor Selection*. NREL/CP-5500-86682. National Renewable Energy Laboratory.

PyTorch Contributors. 2026a. *Distributed Checkpoint: Torch.distributed.checkpoint*.

PyTorch Contributors. 2026b. *Torch Distributed Elastic Rendezvous*. PyTorch documentation.

PyTorch Contributors. 2026c. *Torchrun (Elastic Launch)*. PyTorch documentation.

Qiao, Aurick, Sang Keun Choe, Suhas Jayaram Subramanya, Willie Neiswanger, Qirong Ho, Hao Zhang, Gregory R. Ganger, and Eric P. Xing. 2021. "Pollux: Co-Adaptive Cluster Scheduling for Goodput-Optimized Deep Learning." *15th USENIX Symposium on Operating Systems Design and Implementation (OSDI 21)*, 1–18.

Quaye, J., A. Parrish, O. Inel, C. Rastogi, H. R. Kirk, M. Kahng, E. Van Liemt, et al. 2024. "Adversarial Nibbler: An Open Red-Teaming Method for Identifying Diverse Harms in Text-to-Image Generation." *The 2024 ACM Conference on Fairness, Accountability, and Transparency*, 388–406. https://doi.org/10.1145/3630106.3658913.

Quiñonero-Candela, Joaquin, Masashi Sugiyama, Anton Schwaighofer, and Neil D. Lawrence, eds. 2009. *Dataset Shift in Machine Learning*. Neural Information Processing Series. The MIT Press.

Radovanovic, Ana, Ross Koningstein, Ian Schneider, Bokan Chen, Alexandre Duarte, Binz Roy, Diyue Xiao, et al. 2021. "Carbon-Aware Computing for Datacenters." *arXiv Preprint arXiv:2106.11750*, ahead of print. https://doi.org/10.48550/arXiv.2106.11750.

Rafailov, Rafael, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. 2023. "Direct Preference Optimization: Your Language Model Is Secretly a Reward Model." *Advances in Neural Information Processing Systems (NeurIPS)* 36: 53728–41. https://doi.org/10.52202/075280-2338.

Raffel, C., N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu. 2020. "Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer." *Journal of Machine Learning Research* 21 (140): 1–67.

Rajbhandari, Samyam, Jeff Rasley, Olatunji Ruwase, and Yuxiong He. 2020. "ZeRO: Memory Optimizations Toward Training Trillion Parameter Models." *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*, 1–16. https://doi.org/10.1109/sc41405.2020.00024.

Raji, Inioluwa Deborah, and Joy Buolamwini. 2019. "Actionable Auditing: Investigating the Impact of Publicly Naming Biased Performance Results of Commercial AI Products." *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society*, 429–35. https://doi.org/10.1145/3306618.3314244.

Ramesh, A., M. Pavlov, G. Goh, S. Gray, C. Voss, A. Radford, M. Chen, and I. Sutskever. 2021. "Zero-Shot Text-to-Image Generation." In *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*, edited by Marina Meila and Tong Zhang, vol. 139. Proceedings of Machine Learning Research. PMLR.

Rashid, Layali, Karthik Pattabiraman, and Sathish Gopalakrishnan. 2015. "Characterizing the Impact of Intermittent Hardware Faults on Programs." *IEEE Transactions on Reliability* 64 (1): 297–310. https://doi.org/10.1109/tr.2014.2363152.

Rasley, Jeff, Samyam Rajbhandari, Olatunji Ruwase, and Yuxiong He. 2020. "DeepSpeed: System Optimizations Enable Training Deep Learning Models with over 100 Billion Parameters." *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 3505–6. https://doi.org/10.1145/3394486.3406703.

Ray Project. 2026. *Ray Train: Handling Failures and Node Preemption*. Ray documentation.

Reagen, B., U. Gupta, L. Pentecost, P. Whatmough, S. K. Lee, N. Mulholland, D. Brooks, and G.-Y. Wei. 2018. "Ares: A Framework for Quantifying the Resilience of Deep Neural Networks." *2018 55th ACM/ESDA/IEEE Design Automation Conference (DAC)*, 1–6. https://doi.org/10.1109/dac.2018.8465834.

Rebuffi, Sylvestre-Alvise, Hakan Bilen, and Andrea Vedaldi. 2017. "Learning Multiple Visual Domains with Residual Adapters." *Advances in Neural Information Processing Systems (NeurIPS)* 30.

Recht, Benjamin, Rebecca Roelofs, Ludwig Schmidt, and Vaishaal Shankar. 2019. "Do ImageNet Classifiers Generalize to ImageNet?" *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 5389–400.

Reddi, Vijay Janapa, Christine Cheng, David Kanter, Peter Mattson, Guenther Schmuelling, Carole-Jean Wu, Brian Anderson, et al. 2019. "MLPerf Inference Benchmark." *2020 ACM/IEEE 47th Annual International Symposium on Computer Architecture (ISCA)*, 446–59. https://doi.org/10.1109/isca45697.2020.00045.

Reddi, Vijay Janapa, and Meeta Sharma Gupta. 2013. *Resilient Architecture Design for Voltage Variation*. Synthesis Lectures on Computer Architecture. Springer International Publishing. https://doi.org/10.1007/978-3-031-01739-1.

Reed, I. S., and G. Solomon. 1960. "Polynomial Codes over Certain Finite Fields." *Journal of the Society for Industrial and Applied Mathematics* 8 (2): 300–304. https://doi.org/10.1137/0108018.

Regenscheid, Andrew. 2018. *Platform Firmware Resiliency Guidelines*. NIST SP 800-193. National Institute of Standards; Technology. https://doi.org/10.6028/NIST.SP.800-193.

Ren, Gang, Eric Tune, Tipp Moseley, Yixin Shi, Silvius Rus, and Robert Hundt. 2010. "Google-Wide Profiling: A Continuous Profiling Infrastructure for Data Centers." *IEEE Micro* 30 (4): 65–79. https://doi.org/10.1109/MM.2010.68.

Reuters. 2023. *Human-Machine Teams Driven by AI Are about to Reshape Warfare*. Reuters news article.

Rhodes, Christopher J. 2019. "Endangered Elements, Critical Raw Materials and Conflict Minerals." *Science Progress* 102 (4): 304–50. https://doi.org/10.1177/0036850419884873.

Ribeiro, M. H., R. Ottoni, R. West, V. A. F. Almeida, and Jr. Meira Wagner. 2020. "Auditing Radicalization Pathways on YouTube." *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 131–41. https://doi.org/10.1145/3351095.3372879.

Ribeiro, Marco Tulio, Sameer Singh, and Carlos Guestrin. 2016. "Why Should I Trust You?: Explaining the Predictions of Any Classifier." *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1135–44. https://doi.org/10.1145/2939672.2939778.

Robertson, J., and M. Riley. 2018. *The Big Hack: How China Used a Tiny Chip to Infiltrate U.S. Companies - Bloomberg*.

Rodio, Angelo, and Giovanni Neglia. 2024. "FedStale: Leveraging Stale Client Updates in Federated Learning." *Frontiers in Artificial Intelligence and Applications (ECAI 2024)* 78. https://doi.org/10.3233/faia240849.

Rolnick, David, Arun Ahuja, Jonathan Schwarz, Timothy Lillicrap, and Greg Wayne. 2019. "Experience Replay for Continual Learning." *Advances in Neural Information Processing Systems (NeurIPS)*.

Rubin, Donald B. 1980. "Randomization Analysis of Experimental Data: The Fisher Randomization Test Comment." *Journal of the American Statistical Association* 75 (371): 591–93. https://doi.org/10.2307/2287650.

Rudin, C. 2019. "Stop Explaining Black Box Machine Learning Models for High Stakes Decisions and Use Interpretable Models Instead." *Nature Machine Intelligence* 1 (5): 206–15. https://doi.org/10.1038/s42256-019-0048-x.

Russell, Jon. 2018. "Strava Says It Will Simplify Privacy Settings and Review App Features After Exposing Military Bases." *TechCrunch*, January.

Russell, S. 2021. "Human-Compatible Artificial Intelligence." In *Human-Like Machine Intelligence*. Oxford University Press. https://doi.org/10.1093/oso/9780198862536.003.0001.

Rusu, A. A., N. C. Rabinowitz, G. Desjardins, H. Soyer, J. Kirkpatrick, K. Kavukcuoglu, R. Pascanu, and R. Hadsell. 2016. "Progressive Neural Networks." *arXiv Preprint arXiv:1606.04671*.

Ryan, R. M., and E. L. Deci. 2000. "Self-Determination Theory and the Facilitation of Intrinsic Motivation, Social Development, and Well-Being." *American Psychologist* 55 (1): 68–78. https://doi.org/10.1037/0003-066x.55.1.68.

Saltzer, J. H., and M. D. Schroeder. 1975. "The Protection of Information in Computer Systems." *Proceedings of the IEEE* 63 (9): 1278–308. https://doi.org/10.1109/proc.1975.9939.

Sangchoolie, Behrooz, Karthik Pattabiraman, and Johan Karlsson. 2017. "One Bit Is (Not) Enough: An Empirical Study of the Impact of Single and Multiple Bit-Flip Errors." *2017 47th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*, 97–108. https://doi.org/10.1109/dsn.2017.30.

Sanh, Victor, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2019. "DistilBERT, a Distilled Version of BERT: Smaller, Faster, Cheaper and Lighter." *arXiv Preprint arXiv:1910.01108*, ahead of print. https://doi.org/10.48550/arXiv.1910.01108.

Schaeffer, Rylan, Brando Miranda, and Sanmi Koyejo. 2023. "Are Emergent Abilities of Large Language Models a Mirage?" *Advances in Neural Information Processing Systems (NeurIPS)*, 55565–81. https://doi.org/10.52202/075280-2425.

Schäfer, Mike S. 2023. "The Notorious GPT: Science Communication in the Age of Artificial Intelligence." *Journal of Science Communication* 22 (02): Y02. https://doi.org/10.22323/2.22020402.

Schmidt, Victor, Kamal Goyal, Aditya Joshi, Boris Feld, Liam Conell, Nikolas Laskaris, Doug Blank, Jonathan Wilson, Sorelle Friedler, and Sasha Luccioni. 2021. *CodeCarbon: Estimate and Track Carbon Emissions from Machine Learning Computing*. Zenodo software citation. https://doi.org/10.5281/zenodo.4658424.

Schmuck, Frank, and Roger Haskin. 2002. "GPFS: A Shared-Disk File System for Large Computing Clusters." *Proceedings of the USENIX Conference on File and Storage Technologies (FAST '02)*, 231–44.

Schneider Electric. 2024. *Schneider Electric Announces New Solutions to Address the Energy and Sustainability Challenges Spurred by AI*.

Schneider, Ian, and Taylor Mattia. 2024. "Carbon Accounting in the Cloud: A Methodology for Allocating Emissions Across Data Center Users." *arXiv Preprint arXiv:2406.09645*.

Schneider, Ian, Hui Xu, Stephan Benecke, David Patterson, Keguo Huang, Parthasarathy Ranganathan, and Cooper Elsworth. 2025. "Life-Cycle Emissions of AI Hardware: A Cradle-to-Grave Approach and Generational Trends." *arXiv Preprint arXiv:2502.01671*.

Schroeder, Bianca, Eduardo Pinheiro, and Wolf-Dietrich Weber. 2009. "DRAM Errors in the Wild: A Large-Scale Field Study." *ACM SIGMETRICS Performance Evaluation Review* 37 (1): 193–204. https://doi.org/10.1145/2492101.1555372.

Schulman, John, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. "Proximal Policy Optimization Algorithms." *arXiv Preprint arXiv:1707.06347*, ahead of print. https://doi.org/10.48550/arXiv.1707.06347.

Schwan, P. 2003. "Lustre: Building a File System for 1,000-Node Clusters." *Proceedings of the 2003 Linux Symposium*, 380–86.

Schwartz, Roy, Jesse Dodge, Noah A. Smith, and Oren Etzioni. 2020. "Green AI." *Communications of the ACM* 63 (12): 54–63. https://doi.org/10.1145/3381831.

Sculley, D., Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-François Crespo, and Dan Dennison. 2015. "Hidden Technical Debt in Machine Learning Systems." *Advances in Neural Information Processing Systems (NeurIPS)* 28: 2503–11.

Seide, Frank, Hao Fu, Jasha Droppo, Gang Li, and Dong Yu. 2014. "1-Bit Stochastic Gradient Descent and Its Application to Data-Parallel Distributed Training of Speech DNNs." *Interspeech 2014*, 1058–62. https://doi.org/10.21437/interspeech.2014-274.

Selvaraju, R. R., M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra. 2017. "Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization." *2017 IEEE International Conference on Computer Vision (ICCV)*, 618–26. https://doi.org/10.1109/iccv.2017.74.

SemiAnalysis. 2026. *InferenceX: Open Source Continuous Inference Benchmark*.

Seong, N. H., D. H. Woo, V. Srinivasan, J. A. Rivers, and H.-H. S. Lee. 2010. "SAFER: Stuck-at-Fault Error Recovery for Memories." *2010 43rd Annual IEEE/ACM International Symposium on Microarchitecture*, 115–24. https://doi.org/10.1109/micro.2010.46.

Sergeev, Alexander, and Mike Del Balso. 2018. "Horovod: Fast and Easy Distributed Deep Learning in TensorFlow." *CoRR* abs/1802.05799.

Sevilla, Jaime, Lennart Heim, Anson Ho, Tamay Besiroglu, Marius Hobbhahn, and Pablo Villalobos. 2022. "Compute Trends Across Three Eras of Machine Learning." *2022 International Joint Conference on Neural Networks (IJCNN)*, 1–8. https://doi.org/10.1109/ijcnn55064.2022.9891914.

Shafahi, Ali, W. Ronny Huang, Mahyar Najibi, Octavian Suciu, Christoph Studer, Tudor Dumitras, and Tom Goldstein. 2018. "Poison Frogs! Targeted Clean-Label Poisoning Attacks on Neural Networks." *Advances in Neural Information Processing Systems* 31.

Shafahi, Ali, Mahyar Najibi, Amin Ghiasi, Zheng Xu, John Dickerson, Christoph Studer, Larry S. Davis, Gavin Taylor, and Tom Goldstein. 2019. "Adversarial Training for Free!" *arXiv Preprint arXiv:1904.12843*, ahead of print. https://doi.org/10.48550/arXiv.1904.12843.

Shalev-Shwartz, S. 2012. "Online Learning and Online Convex Optimization." *Foundations and Trends in Machine Learning* 4 (2): 107–94. https://doi.org/10.1561/2200000018.

Shalev-Shwartz, Shai, Shaked Shammah, and Amnon Shashua. 2017. "On a Formal Model of Safe and Scalable Self-Driving Cars." *ArXiv Preprint* abs/1708.06374.

Shallue, Christopher J, Jaehoon Lee, Joseph Antognini, Jascha Sohl-Dickstein, Roy Frostig, and George E Dahl. 2019. "Measuring the Effects of Data Parallelism on Neural Network Training." *Journal of Machine Learning Research* 20: 1–49.

Shan, S., W. Ding, J. Passananti, S. Wu, H. Zheng, and B. Y. Zhao. 2023. "Nightshade: Prompt-Specific Poisoning Attacks on Text-to-Image Generative Models." *2024 IEEE Symposium on Security and Privacy (SP)* abs/2310.13828: 807–25. https://doi.org/10.1109/sp54263.2024.00207.

Shazeer, Noam, Azalia Mirhoseini, Krzysztof Maziarz, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. 2017. "Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer." *International Conference on Learning Representations (ICLR)*.

Sheaffer, J. W., D. P. Luebke, and K. Skadron. 2007. "A Hardware Redundancy and Recovery Mechanism for Reliable Scientific Computation on Graphics Processors." *Graphics Hardware* 2007: 55–64. https://doi.org/10.2312/eggh/eggh07/055-064.

Shehabi, Arman, Sarah Josephine Smith, Alex Hubbard, Alexander Newkirk, Nuoa Lei, Md Abu Bakar Siddik, Billie Holecek, Jonathan G. Koomey, Eric R. Masanet, and Dale A. Sartor. 2024. *2024 United States Data Center Energy Usage Report*. LBNL-2001637. Lawrence Berkeley National Laboratory. https://doi.org/10.71468/P1WC7Q.

Sheng, Ying, Shiyi Cao, Dacheng Li, Coleman Hooper, Shuo Chen, Daxin Jiang, Christopher Rodriguez, et al. 2024. "S-LoRA: Serving Thousands of Concurrent LoRA Adapters." *Proceedings of the 7th Conference on Machine Learning and Systems (MLSys)*.

Shewhart, Walter A. 1931. *Economic Control of Quality of Manufactured Product*. D. Van Nostrand Company.

Shneiderman, B. 2020. "Bridging the Gap Between Ethics and Practice: Guidelines for Reliable, Safe, and Trustworthy Human-Centered AI Systems." *ACM Transactions on Interactive Intelligent Systems* 10 (4): 1–31. https://doi.org/10.1145/3419764.

Shneiderman, B. 2022. *Human-Centered AI*. Oxford University Press. https://doi.org/10.1093/oso/9780192845290.001.0001.

Shoeybi, Mohammad, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. 2019. "Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism." *arXiv Preprint arXiv:1909.08053*.

Shokri, Reza, Marco Stronati, Congzheng Song, and Vitaly Shmatikov. 2017. "Membership Inference Attacks Against Machine Learning Models." *2017 IEEE Symposium on Security and Privacy (SP)*, 3–18. https://doi.org/10.1109/SP.2017.41.

Shumailov, Ilia, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. 2024. "AI Models Collapse When Trained on Recursively Generated Data." *Nature* 631 (8022): 755–59. https://doi.org/10.1038/s41586-024-07566-y.

Sigelman, Benjamin H, Luiz André Barroso, Mike Burrows, Pat Stephenson, Manoj Plakal, Donald Beaver, Saul Jaspan, and Chandan Shanbhag. 2010. *Dapper, a Large-Scale Distributed Systems Tracing Infrastructure*. Dapper-2010-1. Google.

Sigler, Eric. 2021. *Scaling Kubernetes to 7,500 Nodes*. OpenAI Blog.

Singh, Narendra, and Oladele A. Ogunseitan. 2022. "Disentangling the Worldwide Web of e-Waste and Climate Change Co-Benefits." *Circular Economy* 1 (2): 100011. https://doi.org/10.1016/j.cec.2022.100011.

Skorobogatov, S. 2009. "Local Heating Attacks on Flash Memory Devices." *2009 IEEE International Workshop on Hardware-Oriented Security and Trust*, 1–6. https://doi.org/10.1109/hst.2009.5225028.

Skorobogatov, S. P., and R. J. Anderson. 2003. "Optical Fault Induction Attacks." In *Cryptographic Hardware and Embedded Systems - CHES 2002*. Springer Berlin Heidelberg. https://doi.org/10.1007/3-540-36400-5_2.

Slade, G. 2007. *Made to Break: Technology and Obsolescence in America*. Harvard University Press. https://doi.org/10.4159/9780674043756.

Smith, Sarah J., Alexander Hubbard, Alex Newkirk, Mohan Ganeshalingam, Billie Holecek, Dale Sartor, Michael Mills, and Arman Shehabi. 2026. *United States Data Center Energy Usage Report: 2025 Update*. LBNL-2001758. Lawrence Berkeley National Laboratory. https://doi.org/10.71468/P1RP4F.

Snell, Charlie, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2025. "Scaling LLM Test-Time Compute Optimally Can Be More Effective Than Scaling Model Parameters." *Proceedings of the 13th International Conference on Learning Representations (ICLR)*.

Song, Liwei, Reza Shokri, and Prateek Mittal. 2019. "Privacy Risks of Securing Machine Learning Models Against Adversarial Examples." *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security*, 241–57. https://doi.org/10.1145/3319535.3354211.

Sridharan, V., N. DeBardeleben, S. Blanchard, K. B. Ferreira, J. Stearley, J. Shalf, and S. Gurumurthi. 2015. "Memory Errors in Modern Systems: The Good, the Bad, and the Ugly." *ACM SIGPLAN Notices* 50 (4): 297–310. https://doi.org/10.1145/2775054.2694348.

Stahel, Walter R. 2016. "The Circular Economy." *Nature* 531 (7595): 435–38. https://doi.org/10.1038/531435a.

Statista. 2022. *Number of Internet of Things (IoT) Connected Devices Worldwide from 2019 to 2030*.

Steck, Harald, Linas Baltrunas, Ehtsham Elahi, Dawen Liang, Yves Raimond, and Justin Basilico. 2021. "Deep Learning for Recommender Systems: A Netflix Case Study." *AI Magazine* 42 (3): 7–18. https://doi.org/10.1609/aimag.v42i3.18140.

Stich, S. U., J.-B. Cordonnier, and M. Jaggi. 2018. "Sparsified SGD with Memory." *Advances in Neural Information Processing Systems* 31.

Stich, Sebastian U. 2019. "Local SGD Converges Fast and Communicates Little." *International Conference on Learning Representations (ICLR)*.

Stoica, Ion, Dawn Song, Raluca Ada Popa, David Patterson, Michael W. Mahoney, Randy Katz, Anthony D. Joseph, et al. 2017. "A Berkeley View of Systems Challenges for AI." *arXiv Preprint arXiv:1712.05855*.

Strubell, Emma, Ananya Ganesh, and Andrew McCallum. 2019. "Energy and Policy Considerations for Deep Learning in NLP." *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, 3645–50. https://doi.org/10.18653/v1/p19-1355.

Subramanya, Suhas Jayaram, Fnu Devvrit, Harsha Vardhan Simhadri, Ravishankar Krishnawamy, and Rohan Kadekodi. 2019. "DiskANN: Fast Accurate Billion-Point Nearest Neighbor Search on a Single Node." *Advances in Neural Information Processing Systems (NeurIPS)* 32.

Sudhakar, Soumya, Vivienne Sze, and Sertac Karaman. 2023. "Data Centers on Wheels: Emissions from Computing Onboard Autonomous Vehicles." *IEEE Micro* 43 (1): 29–39. https://doi.org/10.1109/mm.2022.3219803.

Sundararajan, Mukund, Ankur Taly, and Qiqi Yan. 2017. "Axiomatic Attribution for Deep Networks." *International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 70: 3319–28.

Sutton, Richard S. 2019. "The Bitter Lesson." *Incompleteideas.net* 43.

Sweeney, Latanya. 2002. "K-ANONYMITY: A MODEL for PROTECTING PRIVACY." *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems* 10 (05): 557–70. https://doi.org/10.1142/s0218488502001648.

Systems, Cerebras. 2021. "Wafer-Scale Deep Learning Acceleration with the Cerebras CS-2." *Cerebras Technical Paper* 2021.

Szegedy, Christian, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. 2013. "Intriguing Properties of Neural Networks." *arXiv Preprint arXiv:1312.6199*.

Tabassi, Elham. 2023. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1. National Institute of Standards; Technology. https://doi.org/10.6028/NIST.AI.100-1.

Tang, Hanlin, Shaoduo Gan, Ammar Ahmad Awan, Samyam Rajbhandari, Conglong Li, Xiangru Lian, Ji Liu, Ce Zhang, and Yuxiong He. 2021. "1-Bit Adam: Communication Efficient Large-Scale Training with Adam's Convergence Speed." *Proceedings of the 38th International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 139: 10118–29.

Taori, Rohan, Achal Dave, Vaishaal Shankar, Nicholas Carlini, Benjamin Recht, and Ludwig Schmidt. 2020. "Measuring Robustness to Natural Distribution Shifts in Image Classification." *arXiv Preprint arXiv:2007.00644*.

Thakur, Rajeev, Rolf Rabenseifner, and William Gropp. 2005. "Optimization of Collective Communication Operations in MPICH." *The International Journal of High Performance Computing Applications* 19 (1): 49–66. https://doi.org/10.1177/1094342005051521.

The Fiber Optic Association. 2020. *Fiber Optic Cable Bend Radius or Diameter*.

The Green Grid. 2007. *Green Grid Data Center Power Efficiency Metrics: PUE and DCIE*. The Green Grid.

The Green Grid. 2011. *Water Usage Effectiveness (WUE): A Green Grid Data Center Sustainability Metric*. White Paper 35. The Green Grid.

The New York Times. 2017. *Google's Self-Driving Car Company Sues Uber over Trade Secrets*. The New York Times.

Thompson, N. C., K. Greenewald, K. Lee, and G. F. Manso. 2021. "Deep Learning's Diminishing Returns: The Cost of Improvement Is Becoming Unsustainable." *IEEE Spectrum* 58 (10): 50–55. https://doi.org/10.1109/mspec.2021.9563954.

Tiwari, Devesh, Saurabh Gupta, James Rogers, Don Maxwell, Paolo Rech, Sudharshan Vazhkudai, Daniel Oliveira, et al. 2015. "Understanding GPU Errors on Large-Scale HPC Systems and the Implications for System Design and Operation." *2015 IEEE 21st International Symposium on High Performance Computer Architecture (HPCA)*, 331–42. https://doi.org/10.1109/hpca.2015.7056044.

Touvron, Hugo, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, et al. 2023. "Llama 2: Open Foundation and Fine-Tuned Chat Models." *arXiv Preprint arXiv:2307.09288*.

Tramèr, Florian, Alexey Kurakin, Nicolas Papernot, Ian Goodfellow, Dan Boneh, and Patrick McDaniel. 2017. "Ensemble Adversarial Training: Attacks and Defenses." *International Conference on Learning Representations* abs/1705.07204.

Tramèr, Florian, Fan Zhang, Ari Juels, Michael K Reiter, and Thomas Ristenpart. 2016. "Stealing Machine Learning Models via Prediction APIs." *25th USENIX Security Symposium (USENIX Security 16)*, 601–18.

Tran, Brandon, Jerry Li, and Aleksander Madry. 2018. "Spectral Signatures in Backdoor Attacks." *Advances in Neural Information Processing Systems (NeurIPS)*.

Trusted Computing Group. 2018. *Hardware Requirements for a Device Identifier Composition Engine*. Family 2.0, Level 00, Revision 78. Trusted Computing Group.

Tsai, Min-Jen, Ping-Yi Lin, and Ming-En Lee. 2023. "Adversarial Attacks on Medical Image Classification." *Cancers* 15 (17): 4228. https://doi.org/10.3390/cancers15174228.

Tsai, Timothy, Siva Kumar Sastry Hari, Michael Sullivan, Oreste Villa, and Stephen W. Keckler. 2021. "NVBitFI: Dynamic Fault Injection for GPUs." *2021 51st Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*, 284–91. https://doi.org/10.1109/dsn48987.2021.00041.

Uddin, Mueen, and Azizah Abdul Rahman. 2012. "Energy Efficiency and Low Carbon Enabler Green IT Framework for Data Centers Considering Green Metrics." *Renewable and Sustainable Energy Reviews* 16 (6): 4078–94. https://doi.org/10.1016/j.rser.2012.03.014.

Ultra Ethernet Consortium. 2025. *Ultra Ethernet Specification V1.0*. Ultra Ethernet Consortium.

UNESCO. 2022. *Recommendation on the Ethics of Artificial Intelligence*. UNESCO.

Uptime Institute. 2022. *Uptime Institute Global Data Center Survey 2022*. Uptime Institute.

Uptime Institute. 2026. *Tier Classification System*. Uptime Institute web page.

U.S. Department of Health and Human Services. 2003. "Summary of the HIPAA Privacy Rule."

U.S. Department of Health and Human Services. 2005. "Summary of the HIPAA Security Rule."

U.S. Department of Health and Human Services. 2026. *The HIPAA Security Rule*. HHS.gov guidance.

Valiant, Leslie G. 1990. "A Bridging Model for Parallel Computation." *Communications of the ACM* 33 (8): 103–11. https://doi.org/10.1145/79173.79181.

Velazco, Raoul, Gilles Foucard, and Paul Peronnard. 2010. "Combining Results of Accelerated Radiation Tests and Fault Injections to Predict the Error Rate of an Application Implemented in SRAM-Based FPGAs." *IEEE Transactions on Nuclear Science* 57 (6): 3500–3505. https://doi.org/10.1109/tns.2010.2087355.

Villalobos, Pablo, Anson Ho, Jaime Sevilla, Tamay Besiroglu, Lennart Heim, and Marius Hobbhahn. 2024. "Position: Will We Run Out of Data? Limits of LLM Scaling Based on Human-Generated Data." *Proceedings of the 41st International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 235: 49523–44.

Vinuesa, Ricardo, Hossein Azizpour, Iolanda Leite, Madeline Balaam, Virginia Dignum, Sami Domisch, Anna Felländer, Simone Daniela Langhans, Max Tegmark, and Francesco Fuso Nerini. 2020. "The Role of Artificial Intelligence in Achieving the Sustainable Development Goals." *Nature Communications* 11 (1): 233. https://doi.org/10.1038/s41467-019-14108-y.

Vogels, Thijs, Sai Praneeth Karimireddy, and Martin Jaggi. 2019. "PowerSGD: Practical Low-Rank Gradient Compression for Distributed Optimization." *Advances in Neural Information Processing Systems (NeurIPS)* 32.

Wachter, Sandra, Brent Mittelstadt, and Chris Russell. 2017. "Counterfactual Explanations Without Opening the Black Box: Automated Decisions and the GDPR." *SSRN Electronic Journal* 31: 841. https://doi.org/10.2139/ssrn.3063289.

Wang, Jianyu, Zachary Charles, Zheng Xu, Gauri Joshi, H. Brendan McMahan, Blaise Aguera y Arcas, Maruan Al-Shedivat, et al. 2021. "A Field Guide to Federated Optimization." *arXiv Preprint arXiv:2107.06917*.

Wang, Shibo, Jinliang Wei, Amit Sabne, Andy Davis, Berkin Ilbeyi, Blake Hechtman, Dehao Chen, et al. 2022. "Overlap Communication with Dependent Computation via Decomposition in Large Deep Learning Models." *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 1*, 93–106. https://doi.org/10.1145/3567955.3567959.

Wang, Y., Q. Yao, J. T. Kwok, and L. M. Ni. 2020. "Generalizing from a Few Examples: A Survey on Few-Shot Learning." *ACM Computing Surveys* 53 (3): 1–34. https://doi.org/10.1145/3386252.

Warden, Pete. 2018. "Speech Commands: A Dataset for Limited-Vocabulary Speech Recognition." *arXiv Preprint arXiv:1804.03209*, ahead of print. https://doi.org/10.48550/arXiv.1804.03209.

Warden, Pete, and Daniel Situnayake. 2020. *TinyML: Machine Learning with TensorFlow Lite on Arduino and Ultra-Low-Power Microcontrollers*. O'Reilly Media.

Waymo. 2026. *Waymo Safety Impact Data Hub*.

Wei, Jason, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, et al. 2022. "Emergent Abilities of Large Language Models." *Transactions on Machine Learning Research* 4. https://doi.org/10.48550/arXiv.2206.07682.

Widmer, Gerhard, and Miroslav Kubat. 1996. "Learning in the Presence of Concept Drift and Hidden Contexts." *Machine Learning* 23 (1): 69–101. https://doi.org/10.1023/a:1018046501280.

Wiener, N. 1960. "Some Moral and Technical Consequences of Automation: As Machines Learn They May Develop Unforeseen Strategies at Rates That Baffle Their Programmers." *Science* 131 (3410): 1355–58. https://doi.org/10.1126/science.131.3410.1355.

Williams, Samuel, Andrew Waterman, and David Patterson. 2009. "Roofline: An Insightful Visual Performance Model for Multicore Architectures." *Communications of the ACM* 52 (4): 65–76. https://doi.org/10.1145/1498765.1498785.

Wongpanich, A., T. Oguntebi, J. Baiocchi Paredes, Y. E. Wang, P. M. Phothilimthana, R. Mitra, Z. Zhou, N. Kumar, and V. J. Reddi. 2025. *Machine Learning Fleet Efficiency: Analyzing and Optimizing Large-Scale Google TPU Systems with ML Productivity Goodput*.

Work of the Future, MIT Task Force on the. 2020. *The Work of the Future: Building Better Jobs in an Age of Intelligent Machines*. Massachusetts Institute of Technology.

World Resources Institute, and World Business Council for Sustainable Development. 2023. *Greenhouse Gas Protocol: Corporate Standard*.

Wright, Peter, and Paul Greengrass. 1987. *Spycatcher: The Candid Autobiography of a Senior Intelligence Officer*. Viking Penguin.

Wu, Carole-Jean, Ramya Raghavendra, Udit Gupta, Bilge Acun, Newsha Ardalani, Kiwan Maeng, Gloria Chang, et al. 2022. "Sustainable AI: Environmental Implications, Challenges and Opportunities." *Proceedings of Machine Learning and Systems* 4: 795–813.

Wulf, Wm. A., and Sally A. McKee. 1995. "Hitting the Memory Wall: Implications of the Obvious." *ACM SIGARCH Computer Architecture News* 23 (1): 20–24. https://doi.org/10.1145/216585.216588.

Wynsberghe, Aimee van. 2021. "Sustainable AI: AI for Sustainability and the Sustainability of AI." *AI and Ethics* 1: 213–18. https://doi.org/10.1007/s43681-021-00043-6.

Xiao, Guangxuan, Ji Lin, Mickaël Seznec, Hao Wu, Julien Demouth, and Song Han. 2023. "SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models." *International Conference on Machine Learning (ICML)*, 38087–99.

Xiao, Wencong, Romil Bhardwaj, Ramachandran Ramjee, Muthian Sivathanu, Nipun Kwatra, Zhenhua Han, Prabal Patel, et al. 2018. "Gandiva: Introspective Cluster Scheduling for Deep Learning." *13th USENIX Symposium on Operating Systems Design and Implementation (OSDI 18)*, 595–610.

Xu, Dianlei, Tong Li, Yong Li, Xiang Su, Sasu Tarkoma, Tao Jiang, Jon Crowcroft, and Pan Hui. 2020. "Edge Intelligence: Architectures, Challenges, and Applications." *arXiv Preprint arXiv:2003.12172*.

Xu, Weilin, David Evans, and Yanjun Qi. 2018. "Feature Squeezing: Detecting Adversarial Examples in Deep Neural Networks." *Proceedings 2018 Network and Distributed System Security Symposium* 26. https://doi.org/10.14722/ndss.2018.23198.

Xu, Yuanzhong, HyoukJoong Lee, Dehao Chen, Blake Hechtman, Yanping Huang, Rahul Joshi, Maxim Krikun, et al. 2021. *GSPMD: General and Scalable Parallelization for ML Computation Graphs*.

Xu, Zheng, Yanxiang Zhang, Galen Andrew, Christopher A. Choquette-Choo, Peter Kairouz, H. Brendan McMahan, Jesse Rosenstock, and Yuanbo Zhang. 2023. "Federated Learning of Gboard Language Models with Differential Privacy." *arXiv Preprint arXiv:2305.18465*.

Yao, Andrew C. 1982. "Protocols for Secure Computations." *23rd Annual Symposium on Foundations of Computer Science (Sfcs 1982)*, 160–64. https://doi.org/10.1109/sfcs.1982.38.

Yin, Dong, Yudong Chen, Kannan Ramchandran, and Peter Bartlett. 2018. "Byzantine-Robust Distributed Learning: Towards Optimal Statistical Rates." *International Conference on Machine Learning (ICML)*, 5650–59.

Yin, Shuang, Xiang Zhou, and Cedric Lam. 2022. *FEC Requirements for 800GbE/1.6TbE Optics: An Operator's Perspective*. IEEE 802.3df Task Force.

Yoo, A. B., M. A. Jette, and M. Grondona. 2003. "SLURM: Simple Linux Utility for Resource Management." In *Job Scheduling Strategies for Parallel Processing*. Springer Berlin Heidelberg. https://doi.org/10.1007/10968987_3.

You, Jie, Jae-Won Chung, and Mosharaf Chowdhury. 2023. "Zeus: Understanding and Optimizing GPU Energy Consumption of DNN Training." *20th USENIX Symposium on Networked Systems Design and Implementation (NSDI 23)*, 119–39.

You, Yang, Igor Gitman, and Boris Ginsburg. 2017. "Large Batch Training of Convolutional Networks." *arXiv Preprint arXiv:1708.03888*.

You, Y., J. Li, S. Reddi, J. Hseu, S. Kumar, S. Bhojanapalli, X. Song, J. Demmel, K. Keutzer, and C.-J. Hsieh. 2020. "Large Batch Optimization for Deep Learning: Training BERT in 76 Minutes." *International Conference on Learning Representations (ICLR)*.

Young, John W. 1974. "A First Order Approximation to the Optimum Checkpoint Interval." *Communications of the ACM* 17 (9): 530–31. https://doi.org/10.1145/361147.361115.

Yu, Gyeong-In, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. 2022. "Orca: A Distributed Serving System for Transformer-Based Generative Models." *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*, 521–38.

Yurdakul, Bilal, and Joshua Naranjo. 2020. "Statistical Properties of the Population Stability Index." *The Journal of Risk Model Validation* 52. https://doi.org/10.21314/jrmv.2020.227.

Zafrir, Ofir, Guy Boudoukh, Peter Izsak, and Moshe Wasserblat. 2019. "Q8BERT: Quantized 8Bit BERT." *2019 Fifth Workshop on Energy Efficient Machine Learning and Cognitive Computing - NeurIPS Edition (EMC2-NIPS)*, 36–39. https://doi.org/10.1109/emc2-nips53020.2019.00016.

Zaharia, Matei, Andrew Chen, Aaron Davidson, Ali Ghodsi, Sue Ann Hong, Andy Konwinski, Siddharth Murching, et al. 2018. "Accelerating the Machine Learning Lifecycle with MLflow." *IEEE Data Engineering Bulletin* 41 (4): 39–45.

Zaharia, M., R. S. Xin, P. Wendell, T. Das, M. Armbrust, A. Dave, X. Meng, et al. 2016. "Apache Spark: A Unified Engine for Big Data Processing." *Communications of the ACM* 59 (11): 56–65. https://doi.org/10.1145/2934664.

Zhang, Jeff, Kartheek Rangineni, Zahra Ghodsi, and Siddharth Garg. 2018. "ThUnderVolt: Enabling Aggressive Voltage Underscaling and Timing Error Resilience for Energy Efficient Deep Learning Accelerators." *2018 55th ACM/ESDA/IEEE Design Automation Conference (DAC)*, 1–6. https://doi.org/10.1109/dac.2018.8465918.

Zhang, Susan, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, et al. 2022. *OPT: Open Pre-Trained Transformer Language Models*.

Zhao, M., and G. E. Suh. 2018. "FPGA-Based Remote Power Side-Channel Attacks." *2018 IEEE Symposium on Security and Privacy (SP)*, 229–44. https://doi.org/10.1109/sp.2018.00049.

Zhao, Yanli, Andrew Gu, Rohan Varma, Liang Luo, Chien-Chin Huang, Min Xu, Less Wright, et al. 2023. "PyTorch FSDP: Experiences on Scaling Fully Sharded Data Parallel." *Proceedings of the VLDB Endowment* 16 (12): 3848–60. https://doi.org/10.14778/3611540.3611569.

Zhao, Yue, Meng Li, Liangzhen Lai, Naveen Suda, Damon Civin, and Vikas Chandra. 2018. "Federated Learning with Non-IID Data." *CoRR* abs/1806.00582.

Zhong, Yinmin, Shengyu Liu, Junda Chen, Jianbo Hu, Yibo Zhu, Xuanzhe Liu, Xin Jin, and Hao Zhang. 2024. "DistServe: Disaggregating Prefill and Decoding for Goodput-Optimized Large Language Model Serving." *18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24)*, July, 193–210.

Zhu, Ligeng, Zhijian Liu, and Song Han. 2019. "Deep Leakage from Gradients." *Advances in Neural Information Processing Systems (NeurIPS)* 32: 14774–84.

Zhu, Y., H. Eran, D. Firestone, C. Guo, M. Lipshteyn, Y. Liron, J. Padhye, S. Raindel, M. H. Yahia, and M. Zhang. 2015. "Congestion Control for Large-Scale RDMA Deployments." *ACM SIGCOMM Computer Communication Review* 45 (4): 523–36. https://doi.org/10.1145/2829988.2787484.

Ziegler, J. F., H. W. Curtis, H. P. Muhlfeld, C. J. Montrose, B. Chin, M. Nicewicz, C. A. Russell, et al. 1996. "IBM Experiments in Soft Fails in Computer Electronics (1978–1994)." *IBM Journal of Research and Development* 40 (1): 3–18. https://doi.org/10.1147/rd.401.0003.

Zou, Andy, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. 2023. *Universal and Transferable Adversarial Attacks on Aligned Language Models*.

Zu, Y., A. Ghaffarkhah, H.-V. Dang, B. Towles, S. Hand, S. Huda, A. Bello, et al. 2024. "Resiliency at Scale: Managing Google's TPUv4 Machine Learning Supercomputer." *21st USENIX Symposium on Networked Systems Design and Implementation (NSDI 24)*, 761–74.

A

Single-Machine Foundations

Every fleet is built from single machines, and this book assumes that the reader already reasons fluently about one of them. This appendix restates that single-machine foundation in reference form so that a reader who meets the prerequisites only loosely can work every calculation in the book without another text. It collects the Data, Algorithm, and Machine (D·A·M) diagnostic framework and the iron law it rests on, arithmetic intensity and the roofline bound, the accelerator memory hierarchy, the kernel-level levers that shorten a single device's step, the byte accounting of mixed-precision training, and the arithmetic of transformer training and inference. The chapters build the fleet-scale Compute, Communication, and Coordination (C^3) framework on this base.

Learning Objectives

- Classify single-machine bottlenecks by dominant Data, Algorithm, or Machine axis while recognizing mixed causes
- Map optimization techniques to their D·A·M intersection zone to understand which axes they span
- Apply the iron law equation to quantitatively diagnose performance problems
- Distinguish between memory-bound and compute-bound workloads using arithmetic intensity and the roofline bound
- Calculate the per-parameter memory of mixed-precision training and the FLOPs of transformer training and decoding
- Select appropriate profiling tools and optimization strategies for each Data, Algorithm, or Machine axis
- Evaluate system health using D·A·M scorecard metrics (I/O overhead, active parameters, MFU)

How to Use This Appendix

This appendix is a reference for single-node performance and for the single-machine quantities that fleet-scale arguments take as given. Start from the symptom or the missing quantity:

- **When a single accelerator is slow**: Use the iron law mapping (section A.2) and the bottleneck diagnostic (section A.3.2) to decide which term binds.
- **When deciding whether a kernel is memory bound or compute bound**: Use the roofline bound (section A.3.1).
- **When a byte must cross a memory boundary**: Use the memory hierarchy (section A.4) to see which tier sets its cost.
- **When a trace shows many short kernels, launch gaps, or attention intermediates written to device memory**: Use the kernel-level levers (section A.5).
- **When sizing model state**: Use the training memory accounting (section A.6) for bytes per parameter.
- **When estimating training compute or decode cost**: Use transformer arithmetic (section A.7).

When training is slow on a single GPU, check utilization, data wait time, and model FLOPs utilization (MFU), then map each to its Data, Algorithm, or Machine axis. When serving misses a latency target, identify whether the regime is latency-bound (overhead), memory-bound (weight/KV movement), or compute-bound. When cost is exploding, use the D·A·M rubric to ensure that effort targets the dominant term, not a nonbottleneck.

The Data · Algorithm · Machine (D·A·M) taxonomy is the diagnostic framework for single-node ML systems engineering: it formalizes how information flow, mathematical logic, and physical execution interact across three diagnostic axes,[1] while recognizing that real systems can involve mixed causes or interactions between axes.

[1] **MECE (mutually exclusive, collectively exhaustive)**: A classification principle from management consulting (popularized by McKinsey) requiring that categories do not overlap and together cover every possibility. Applied to systems engineering, it is useful as an idealized decomposition, but D·A·M bottlenecks often overlap in practice. A kernel can be simultaneously memory-bound, synchronization-heavy, and poorly matched to available hardware.

A.1 Diagnostic Summary

When performance stalls or behavior degrades on a single node, the first diagnostic question is whether the flow of work is blocked by memory movement, arithmetic volume, or hardware throughput. The D·A·M taxonomy maps these three failure modes directly to the iron law of ML systems, the single-machine equation that section 1.5.4 restates before extending it to the Fleet Law. Each axis plays a distinct role, binds on a distinct physical quantity (bandwidth, operations, or peak throughput), and answers to its own high-leverage optimization (table A.1).

Table A.1: D·A·M Axis Reference: Each axis maps to a distinct physical constraint and a high-impact optimization strategy. Start diagnosis here: identify which constraint is binding, then follow the optimization lever.

Axis	Role	Physical constraint	High-leverage optimization
Data (D)	**Information** (The Fuel)	Bandwidth (BW)	I/O Pipeline Optimization
Algorithm (A)	**Logic** (The Blueprint)	Operations (O)	Model Compression
Machine (M)	**Physics** (The Engine)	Throughput (R_{peak})	Hardware Acceleration

The three axes seldom bind in isolation. Pruning weights to cut operations can fragment memory access, and a denser storage format that saves bytes can saturate the host CPUs that must decode it before the accelerator sees any data. Because an intervention on one axis shifts pressure onto the other two, diagnosis has to account for the regions where axes overlap. Figure A.1 draws the three axes as overlapping sets. Each circle carries the factor its axis contributes to equation A.1, and each pairwise overlap names the question its two axes must answer together. Data and Algorithm decide what the model learns from, Data and Machine decide how information moves, and Algorithm and Machine decide how computation executes efficiently. A bottleneck that sits in an overlap needs a fix that respects both of its axes, and the center, where all three meet, is the domain of ML systems engineering.

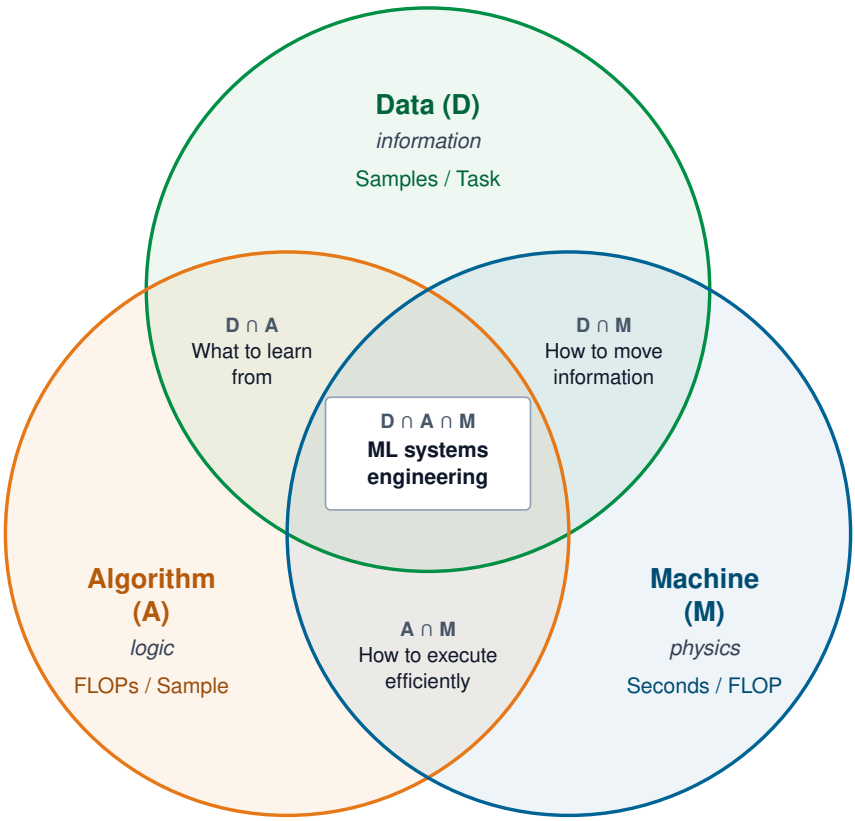

Figure A.1: The D·A·M Intersection Landscape: Each circle is one diagnostic axis (Data, information; Algorithm, logic; Machine, physics), labeled with the factor it contributes to execution time in equation A.1. Each pairwise overlap names the question its two axes must answer jointly (what to learn from, how to move information, and how to execute efficiently). At the center, where all three meet, is ML systems engineering.

A.2 Iron Law Mapping

When a training step or serving request runs slower than expected on a single GPU, the first question is whether the time went to moving bytes, executing arithmetic, or waiting on dispatch overhead. The iron law answers that question in two steps. The first step factors execution time along the three D·A·M axes. A task processes some number of samples (examples in vision or tabular workloads, tokens in language models), each sample costs some number of operations, and the machine takes some time to complete each operation. By dimensional cancellation, total execution time T (seconds) is the product of these three factors (equation A.1):

$$T = \underbrace{\left(\frac{\text{Samples}}{\text{Task}}\right)}_{\text{Data (D)}} \times \underbrace{\left(\frac{\text{FLOPs}}{\text{Sample}}\right)}_{\text{Algorithm (A)}} \times \underbrace{\left(\frac{\text{Seconds}}{\text{FLOP}}\right)}_{\text{Machine (M)}} \tag{A.1}$$

The first two factors fix the total arithmetic work O. Training a model with P parameters on D tokens costs roughly $6PD$ operations across the forward and backward passes (equation A.4), and inference costs roughly $2P$ operations per token in the forward pass. The third factor is the realized time per operation, and no real machine achieves its datasheet value of $1/R_{\text{peak}}$.

The second step explains the gap. Executing an operation requires moving its operands through memory and interconnect, driving the arithmetic units at some realized utilization, and paying fixed control costs. Assuming these three costs occur in sequence, the iron law unpacks execution time into their sum (equation A.2):

$$T = \underbrace{\frac{D_{\text{vol}}}{\text{BW}}}_{\text{Memory time}} + \underbrace{\frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}}_{\text{Compute time}} + \underbrace{L_{\text{lat}}}_{\text{Overhead}} \tag{A.2}$$

Data enters the two equations in different units. In equation A.1, the Data factor counts logical samples or tokens. In equation A.2, D_{vol} counts the physical bytes that cross a memory or interconnect bus, which include model weights, activations, cached sequence state, and gradients as well as the inputs themselves. Batch-1 decode of a language model shows how far apart the two can be. The task processes a single token, yet D_{vol} spans the gigabytes of weights streamed from high-bandwidth memory to produce it.

Algorithm and Machine share the compute term and are separated by which variable the engineer controls. Reducing the total operations (O) is an algorithm lever, while improving the hardware's peak throughput (R_{peak}) or utilization (η_{hw}) is a machine lever. The memory term belongs to no single axis. The Data axis sets how many inputs move, the Algorithm axis sets how many weights and activations move with them, and the Machine axis supplies the bandwidth BW that carries them.

A.2.1 D·A·M coordination from sum to max

The additive iron law represents **sequential execution**, the worst case where Data, Algorithm, and Machine take turns. Overlap can move the sum toward a max, with perfect overlap defining the ideal lower bound:

$$T_{\text{sequential}} = \frac{D_{\text{vol}}}{\text{BW}} + \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}} + L_{\text{lat}} \quad \xrightarrow{\text{overlap}} \quad T_{\text{pipelined}} \geq \max\left(\frac{D_{\text{vol}}}{\text{BW}}, \frac{O}{R_{\text{peak}} \cdot \eta_{\text{hw}}}\right) + L_{\text{lat}}$$

The systems engineer's job is to make these components run in parallel, not in series. Prefetching, separate CUDA streams, asynchronous gradient synchronization, and double buffering each let one axis proceed while another is in flight (table A.2).

Table A.2: D·A·M Overlap Techniques: Each technique allows one D·A·M axis to execute while another is in flight, converting the iron law's additive terms into overlapped terms.

Technique	D·A·M axes overlapped	Implementation
Prefetching	D overlaps M	DataLoader with `prefetch_factor`, `pin_memory=True`
CUDA Streams	D overlaps M	Separate streams for H2D transfer and compute
Async Gradient Sync	M (communication) overlaps A	Overlap gradient AllReduce with remaining backward computation as gradient buckets become ready
Double Buffering	D overlaps M	Fill buffer $i+1$ while computing on buffer i

A.3 Arithmetic Intensity Boundary

This overlap view still leaves one practical question: whether a workload is waiting on bytes or on FLOPs. The boundary between Data (memory bound) and Machine (compute bound) is not arbitrary; it is defined mathematically by the arithmetic intensity[2] (I) of the workload.

The symptoms below are screening signals, not diagnoses. A workload whose arithmetic intensity sits left of $R_{\text{peak}}/\text{BW}$ cannot keep the accelerator fed with enough useful work per byte. Raw utilization must still be read alongside memory-bandwidth utilization, data-loader wait, kernel occupancy, and trace gaps.

- **Low GPU utilization** ($<$ 80 percent): Screen for data, CPU, or launch starvation. Confirm with data-loader wait, host CPU saturation, memory-bandwidth counters, and timeline gaps.
- **High GPU utilization** ($>$ 95 percent): Treat this as machine bound only if compute units are busy and memory bandwidth is not saturated. If memory bandwidth is saturated, the bottleneck is still on the Data/Machine boundary.
- **If batch size is 1**: Treat this as a warning that latency or launch overhead may matter. For large language model (LLM) decode, batch-1 execution is often memory-bandwidth bound unless traces show dispatch gaps.
- **If arithmetic intensity is below the ridge point** (about 295.2 FLOP/byte on H100): The workload is likely memory bound (Data/Machine boundary).

[2] **Arithmetic intensity**: The ratio of floating-point operations to bytes transferred (FLOP/byte). It determines whether a workload is memory-bound or compute-bound by comparison against the hardware's *ridge point* ($R_{\text{peak}}/\text{BW}$).

A.3.1 The roofline bound

Arithmetic intensity is a property of the workload and its implementation, not of the hardware. A large matrix multiply reuses each weight it loads across many multiply-accumulates and so has high intensity. An element-wise operation, or the generation of a single token, uses each loaded value only once or twice. The ridge point belongs to a simple upper bound on single-device throughput, the Roofline Model (Williams et al. 2009). A kernel with arithmetic intensity I, running on a device with peak compute rate R_{peak} and memory bandwidth BW, can attain at most

$$R_{\text{attain}}(I) = \min\left(R_{\text{peak}},\ I \times \text{BW}\right) \tag{A.3}$$

On log-log axes, the bound is a bandwidth roof of slope one ($I \times \text{BW}$) that meets a flat compute roof (R_{peak}) at the ridge point $I_{\text{ridge}} = R_{\text{peak}}/\text{BW}$. Left of the ridge, a kernel is memory bound, and only more bandwidth or fewer bytes per operation raises its ceiling. Right of the ridge, it is compute bound, and only a faster or better-used datapath helps. In D·A·M terms, the bandwidth roof belongs to the Data axis, the compute roof to the Machine axis, and the Algorithm axis decides where on the intensity axis the kernel sits. The roofline is a ceiling, not a prediction. A kernel far below both roofs is limited by something the model omits, such as launch overhead, synchronization, or low occupancy. Figure 2.8 draws both roofs for an H100 and places decode, prefill, and training workloads against them.

For an H100, the book's hardware constants give a dense FP16 peak of 989 TFLOP/s and a high-bandwidth memory (HBM) bandwidth of 3.35 TB/s, which places the ridge at about 295.2 FLOP/byte. Batch-1 decode of a 7-billion-parameter language model stored in 16-bit precision performs about two operations per parameter for each generated token (one multiply and one add). It must read every two-byte weight once, so its intensity is about 1 FLOP/byte. Equation A.3 then

caps attainable throughput at 3.35 TFLOP/s, roughly 0.3 percent of peak, however well the kernels are written. Serving B sequences in one decode step reuses each loaded weight B times and raises the intensity to roughly B FLOP/byte, until key-value cache traffic, which grows with batch size and context length, begins to dominate the bytes moved. That arithmetic is why batching is the first lever of every serving system. Section C.3.3 carries the same reasoning to the network, where interconnect bandwidth takes the place of HBM bandwidth.

A.3.2 Bottleneck diagnostic

Each optimization shortens only one term of the iron law, so adding peak FLOP/s to a memory-bound kernel or adding memory bandwidth to a compute-bound kernel leaves step time unchanged while wasting engineering effort (table A.3).

Table A.3: What Works vs. What Is Wasted: Optimizing the wrong term yields limited improvement and poor cost efficiency. A memory-bound model will not speed up from more peak FLOP/s alone; it needs higher arithmetic intensity, less data movement, or a faster memory subsystem.

If the workload is…	Dominant term	Optimization that works	Optimization that is wasted
Memory-Bound	D_{vol}/BW	Quantization, pruning, batching, kernel fusion, higher memory bandwidth	More peak FLOP/s alone
Compute-Bound	$O/(R_{\text{peak}} \cdot \eta_{\text{hw}})$	Better kernels, Tensor Cores, faster GPU, lower precision	More memory bandwidth (not binding)
Latency-Bound	L_{lat}	Batching requests, kernel fusion, async dispatch	Neither compute nor bandwidth (overhead dominates)

A.4 Accelerator Memory Hierarchy

An accelerator is a fast arithmetic engine surrounded by memories that trade bandwidth for capacity, and most performance arguments in this book reduce to the question of which boundary a byte must cross. Nearest the compute units, registers and on-chip shared memory (SRAM) hold the operands of the kernel that is executing, backed by a shared on-chip L2 cache (50 MB on an H100). Beyond the die, HBM holds the model's working state (80 GB at 3.35 TB/s on an H100). Host memory is reached over PCIe at 64 GB/s per direction, roughly 52× less than HBM, and local NVMe storage sits below host memory.

A kernel whose working set stays on chip runs near arithmetic speed; one that streams operands from HBM is bounded by the roofline above; one that must reach host memory or storage is slower still. The **memory wall**, the widening gap between arithmetic throughput and the bandwidth that feeds it, is the single-machine form of the constraint that reappears at fleet scale as the network wall. A multi-accelerator node adds a scale-up fabric between devices (NVLink, at 450 GB/s per direction on an H100), and the scale-out network adds tiers beyond the node. Section C.2 tabulates those tiers.

A.5 Kernel-Level Levers

Four techniques shorten one device's step without changing the model's mathematics. Each removes a different cost in the iron law of section A.2, and each counts as a gain only when a measured trace confirms it. The four differ in which term they remove, what limits them, and what evidence a trace shows (table A.4).

A.5.1 Fusion

Operator fusion combines consecutive operations into one kernel, so an intermediate result stays in registers or shared memory instead of being written to HBM and read back by the next kernel. It shrinks the data volume D_{vol} of the data term and removes one launch per fused operation from the overhead term L_{lat}. The common kinds are chains of element-wise operations, an element-wise operation fused with a reduction such as a normalization, and operator-specific fusions such as a matrix multiply that applies its bias and activation before writing its output. Fusion stops paying when the live values of the combined kernel exceed the registers and shared memory available, which lowers occupancy or spills to memory, and when a reduction needs results from every thread block. A trace confirms it with fewer kernels and fewer HBM bytes before the step time falls.

A.5.2 Tiling and IO-aware attention

Tiling sizes blocks of a computation to fit on-chip memory, so each operand fetched from HBM feeds many operations and arithmetic intensity rises. IO-aware attention such as FlashAttention applies the idea to attention. It computes the scores block by block, keeping a running maximum and sum for the softmax, so the full matrix of scores between every pair of positions is never written to HBM. The workspace grows linearly with sequence length rather than quadratically. The price is extra rescaling arithmetic and, in training, recomputing the scores in the backward pass. Decode processes one query row per step and has no quadratic intermediate, so its fused attention kernel is a single pass over the KV cache. Tiling reduces D_{vol} and frees HBM capacity for the KV cache and larger batches.

A.5.3 Graph replay

Graph replay, offered on NVIDIA GPUs as CUDA Graphs, records a stable sequence of kernels and memory copies once and then submits the whole sequence with one request in later steps. It removes the host-dispatch part of L_{lat}, which matters when many short kernels leave launch gaps between them, as in batch-one decode. A captured graph replays only the shapes, memory addresses, and control flow it recorded, so a server keeps one graph per shape bucket, and decode captures more easily than prefill, whose shapes vary with prompt length. Replay fuses nothing. The same kernels run with the gaps between them removed, which a timeline shows directly.

A.5.4 Compilation

Graph compilation automates the three levers above. A compiler captures the program as a graph, rewrites and decomposes its operations, fuses them, selects or generates a kernel for each fused group (often by timing candidates, which is called autotuning), and plans memory. Just-in-time compilers keep the eager programming model, ahead-of-time inference engines specialize the program to a declared range of shapes, and tile-level kernel languages let an engineer write a fusion the compiler missed. Compilation reduces both D_{vol} and L_{lat}, at the cost of build and warmup time, regions split at graph breaks where the compiler cannot capture the program, and recompilation when a new shape arrives. Section 9.9.3 prices that cost across a fleet and finds when it pays back.

Table A.4: Kernel-Level Levers: Each lever shortens a different term of one device's step, so a trace must show that term moving before a step-time gain counts.

Lever	Removes	Iron-law term	Main limit	Trace evidence
Fusion	HBM round-trips of intermediates; kernel launches	Data volume; overhead	Register and shared-memory pressure	Fewer kernels and fewer HBM bytes
Tiling and IO-aware attention	Repeated HBM fetches; materialized attention scores	Data volume	Rescaling arithmetic; recomputation	Fewer HBM bytes; smaller attention memory
Graph replay	Host dispatch of each kernel	Overhead	Fixed shapes, addresses, and control flow	Launch gaps disappear from the timeline
Compilation	All three, applied automatically	Data volume and overhead	Build time; graph breaks; recompilation	Fewer, fused kernels after warmup

Introduction to Machine Learning Systems develops these mechanisms in full, for a reader who wants more than this recap. Here, the four rows of table A.4 serve as the checklist for verifying whether a local optimization has moved the intended iron-law term.

A.6 Numeric Precision and Training Memory

Whether a large model fits in an 80 GB H100 memory bank or must be partitioned across dozens of accelerators depends on how many bytes each parameter and optimizer state occupies. In the number formats listed in the notation conventions at the front of the book, FP32 occupies 4 bytes, BF16 and FP16 occupy 2, and FP8 and INT8 occupy one. **Mixed-precision training** runs the forward

and backward passes in a 16-bit format for speed and memory, while it keeps an FP32 master copy of the weights so that small updates are not lost to rounding. The Adaptive Moment Estimation (Adam) optimizer adds two FP32 moment estimates per parameter. Table A.5 totals the resulting bytes per parameter.

Table A.5: Mixed-Precision Adam Memory per Parameter: Persistent training state before activations, framework metadata, and communication buffers. The checkpoint omits gradients because a resumed run recomputes them in its first step.

State	Format	Bytes per parameter	In a resumable checkpoint
Working weights	BF16	2	Yes
Gradients	BF16	2	No (recomputed on resume)
Master weights	FP32	4	Yes
Adam moments (two)	FP32	8	Yes
Total		16 (training state)	14 bytes

For Archetype A, the 175B-parameter dense model the chapters use as their running instance, the training state is 2.8 TB and a resumable checkpoint is 2.45 TB. A single H100 offers 80 GB of HBM (a binary nameplate, about 85.9 GB in decimal units), so the persistent state alone requires at least 33 devices before a single activation is stored. That floor is why the model state must be partitioned across devices, which chapter 6 develops. Serving the same model needs only the 16-bit weights, 350 GB, plus the key-value cache.

Activations are the other large consumer of training memory. The intermediate tensors that the backward pass needs grow with batch size, sequence length, hidden width, and depth, and for long sequences they can exceed the parameter state. **Activation checkpointing** stores only selected activations and recomputes the rest during the backward pass, trading additional compute for memory.

A.7 Transformer Arithmetic

Most worked examples in this book use decoder-only transformers, and three counting rules cover nearly all of their arithmetic. The symbols follow the book's deep learning notation, with P parameters, D training tokens, N_L layers, hidden dimension d, and sequence length S.

Each transformer layer holds four $d \times d$ attention projections and a feed-forward block that expands to $4d$ and back, about $12d^2$ weights in total. A model thus has roughly $P \approx 12 N_L d^2$ parameters excluding embeddings. GPT-3 has $N_L = 96$ and $d = 12{,}288$, which gives 173.9B against its published 175B.

Every weight contributes one multiply and one add per token, so a forward pass costs about $2P$ FLOPs per token. The backward pass computes gradients with respect to both activations and weights and costs about twice the forward pass, so training costs about $6P$ FLOPs per token and

$$O_{\text{train}} \approx 6PD \tag{A.4}$$

in total. For GPT-3, trained on 300 billion tokens, equation A.4 gives 3.15×10^{23} FLOPs, within 1 percent of the reported 3.14×10^{23} FLOPs. The rule omits the attention-score products, which add at most about $S/(6d)$ to the weight FLOPs, or 2.8 percent for GPT-3 at $S = 2{,}048$. That fraction stays small while the context is short compared with the hidden width and grows with long contexts.

Serving an autoregressive model has two phases with different physics. *Prefill* processes the whole prompt at once as matrix-matrix products, so its arithmetic intensity is high and it tends to be compute bound. *Decode* generates one token per step for each sequence as matrix-vector products; at batch size 1 it reads every weight to perform two FLOPs per weight, the memory-bound regime that the roofline example above quantifies. To avoid recomputing attention over the whole prefix at every step, decode stores each layer's keys and values in a **key-value (KV) cache** of $2 N_L d\, s_{\text{elem}}$ bytes per token for standard multihead attention, or 4.72 MB per token for GPT-3 in FP16. Section F.1.2 sizes that cache against device memory.

A.8 Tooling Map

Back-of-envelope byte and FLOP counts narrow the likely bottleneck, but hardware and runtime counters must confirm which iron-law term actually binds before an engineer rewrites a kernel or data pipeline. Each D·A·M axis is exposed by a distinct set of profiling utilities and primary metrics (table A.6).

Table A.6: D·A·M Tooling Map: Profiling utilities for diagnosing bottlenecks along each D·A·M axis.

Axis	Key metric	Primary tool	Secondary tool
Data	Batch Load and Wait Time	Framework profiler trace	`iostat`, `pidstat` (I/O/CPU)
Algorithm	FLOPs, Model Depth	PyTorch Profiler	DeepSpeed Flops Profiler
Machine	GPU Utilization, SM Occupancy	Nsight Compute	DCGM, Nsight Systems

A.9 D·A·M Scorecard

The efficiency ratios in table A.7 grade a system's performance against its theoretical limit. This report card is anchored by MFU,[3] the ratio of achieved model FLOP/s to the hardware's theoretical peak FLOP/s. The MFU grades are illustrative targets for compute-dominated training; memory-bound or latency-bound workloads must be judged against their attainable roofline and service-level objective (SLO) latency.

[3] **MFU (model FLOPs utilization)**: Counts only *useful* model computation in the numerator. Gradient synchronization, memory management, and other overhead within the measurement window remain in elapsed time and lower the reported ratio.

Table A.7: The D·A·M Efficiency Rubric: Use these three numbers to characterize single-machine maturity.

Axis	Metric	Definition	Failing grade	Passing grade
Data	**I/O Overhead**	$\frac{\text{Data Wait Time}}{\text{Total Step Time}}$	$> 10\%$	$< 1\%$
Algorithm	**Compression/Sparsity Ratio**	$\frac{\text{Effective or active parameters}}{\text{Dense baseline parameters}}$	Workload-dependent	Workload-dependent
Machine	**MFU**	$\frac{\text{achieved model FLOP/s}}{\text{peak FLOP/s}}$	$< 30\%$	$> 50\%$

A.10 Scaling the Taxonomy from Node to Fleet

These single-node grades become most useful when they are carried forward into distributed design. The D·A·M taxonomy is the diagnostic baseline for a single machine, but as systems scale from a single node to the machine learning fleet, each axis undergoes a qualitative transformation.

The evolution of constraints

Crossing the node boundary replaces local bus limits and single-device failure rates with cluster-wide fabric and power limits. Scaling from one node to a fleet changes the kind of constraint on every axis, not only its size (table A.8).

Table A.8: D·A·M Constraints from Node to Fleet: How each axis of the D·A·M taxonomy shifts when scaling from a single node to a multi-node fleet. The node-level focus (left) gives way to a qualitatively different fleet-scale constraint (right); the bottleneck migrates from local resource to global coordination cost.

Axis	Node-level focus (D·A·M)	Fleet-scale transformation
Data (D)	I/O Bandwidth (Disk/PCIe)	**The communication wall**: the bottleneck shifts from local storage to the bisection bandwidth of the network fabric.
Algorithm (A)	Model Depth/Ops Count	**The parallelism strategy**: the logic now includes how computation is partitioned across N devices (3D parallelism).
Machine (M)	Peak TFLOP/s, HBM	**The power and reliability wall**: the constraint is no longer just silicon speed, but watts per rack and cluster-wide MTBF.

Bridging to C^3

While D·A·M diagnoses the *components* of a single node, the C^3 taxonomy (Appendix B) diagnoses the *interactions* of the fleet. Each of its three dimensions inherits a specific piece of the single-machine diagnostic model:

1. **Compute (C_1)** inherits the Algorithm and Machine axes, now evaluated per device after the work is partitioned across N devices.
2. **Communication (C_2)** inherits the Data axis, but is governed by network bandwidth, topology, and, over long distances, the speed of light rather than by local I/O alone.
3. **Coordination (C_3)** makes coordination costs explicit at fleet scale. Synchronization, checkpoints, and failure recovery also occur on one node, but distribution expands their scope and couples them to other machines.

This progression exposes cases where higher single-node MFU comes at the expense of scaling efficiency, so the complete fleet outcome can determine the choice. The node is optimized to serve the fleet.

A.11 Summary

The D·A·M taxonomy provides the diagnostic baseline for every ML systems analysis. By isolating the bottleneck to Data, Algorithm, or Machine, practitioners ensure that optimization efforts target the binding constraint. The memory hierarchy, the byte accounting of mixed-precision training, and the transformer counting rules supply the quantities that fleet-scale arguments take as given. This single-node discipline is the prerequisite for the fleet-scale engineering addressed throughout this book.

Key Takeaways: Single-machine diagnostic heuristics

- **Identify the dominant axis**: Decide whether Data, Algorithm, or Machine is the binding constraint before proposing any optimization.
- **Profile arithmetic intensity**: Arithmetic intensity, read against the roofline's ridge point, quantitatively distinguishes between Data-bound and Machine-bound regimes.
- **Use the iron law**: The iron law transforms vague symptoms into specific term-based bottlenecks.
- **Count bytes and FLOPs per parameter**: Mixed-precision Adam training holds 16 bytes per parameter (14 in a checkpoint), and dense transformer training costs about $6PD$ FLOPs. These two rules size most memory and compute budgets in the chapters.
- **Grade with the scorecard**: The Data-Algorithm-Machine scorecard standardizes "good" performance before moving to fleet scale.

B

The C^3 Taxonomy

In a distributed training cluster, "it is slow" is even less informative than on a single machine. A job can stall because accelerators are underutilized (Compute), gradient synchronization saturates interconnects (Communication), or checkpointing and recovery delay progress (Coordination). This appendix extends single-node diagnostics to fleet-scale infrastructure, providing the C^3 scorecard used throughout the book to isolate the binding bottleneck across distributed systems.

How to Use This Appendix

This appendix is designed as a symptom-first reference for fleet-scale diagnosis. Start with the observed failure mode, form a hypothesis about which C^3 axis dominates, and select the measurement that confirms or falsifies that hypothesis:

- **When training throughput is low**: Check model FLOPs utilization (MFU), communication fraction, and runtime goodput in the diagnostic summary (section B.2) and traffic light thresholds (section B.5.1) to isolate the compute, communication, or coordination axis.
- **When scaling efficiency drops as GPUs are added**: Use the Fleet Law decomposition (section B.3) and the intersection landscape (section B.4) to identify which noncompute term grew.
- **When a production job stalls, restarts, or fires hardware alerts**: Match the symptom in the troubleshooting matrix (section B.7), select a profiler from the tooling map (section B.8), or trace silicon counters through the telemetry decision tree (section B.10).
- **When translating a scaling-law compute budget into cluster capacity**: Use the C^3 scorecard (section B.9) and the effective FLOP/s model (section B.11) so optimization effort targets the dominant loss term rather than a nonbottleneck.

B.1 From D·A·M to C^3

Where the single-machine D·A·M framework (Appendix A) diagnoses bottlenecks within one node (data starvation, algorithmic overhead, or hardware saturation), a distributed job can stall even when every individual node is healthy. The **C^3 taxonomy** is the diagnostic framework for fleet-scale ML systems engineering. It isolates the dominant cluster-level axis across Compute (whether the accelerators are doing useful math), Communication (whether the network is moving data fast enough), and Coordination (whether the system is spending too much time on synchronization, failure recovery, and scheduling), along with the intersections between them. Rather than replacing D·A·M, C^3 extends each single-machine axis to the new failure modes that appear when a workload spans a fleet (table B.1).

Table B.1: Mapping D·A·M to C^3: How single-machine diagnostic concerns generalize to fleet-scale challenges.

D·A·M axis	Single-machine concern	C^3 extension	What changes at fleet scale
Data	Host disk → GPU memory pipeline	Communication	Data movement crosses network fabrics; AllReduce, AllGather, and pipeline transfers dominate

D·A·M axis	Single-machine concern	C^3 extension	What changes at fleet scale
Algorithm	Operator choice, autograd graph	Compute/Coordination	Workload must be partitioned (N-way sharding); communication-to-compute ratio determines scaling
Machine	SM occupancy, HBM bandwidth	Coordination	Single-device faults become cluster-wide stragglers; failure recovery, scheduling, and power density constrain N

The most important row in table B.1 is the last one. On a single machine, the overhead term (L_{lat}) in the iron law is typically small, consisting of kernel launch latency, Python dispatch, and synchronization barriers. At fleet scale, coordination becomes an axis in its own right: checkpoint writes, failure detection and recovery, pipeline bubble overhead, scheduler preemptions, and maintenance windows collectively consume a significant fraction of wall time. Coordination is the axis that single-machine diagnosis lacks, and much of this book is devoted to it.

B.2 Diagnostic Summary

When a distributed training job runs below expected throughput, the symptom appears first as idle tensor cores, exposed collective transfers, or frequent barrier and restart stalls. Matching that symptom to the C^3 axis that constrains the fleet links each physical bottleneck directly to its headline metric and highest-leverage optimization (table B.2).

Table B.2: C^3 Diagnostic Summary: Each axis maps to a distinct physical constraint and a high-impact optimization strategy. Start diagnosis here: identify which constraint binds, then follow the optimization pointer to the relevant chapter.

C^3 axis	Physical constraint	Symptoms	Key metric	High-leverage optimization
Compute (C_1)	Arithmetic throughput ($R_{\text{peak}} \times \eta_{\text{hw}}$)	Low MFU, GPU utilization below 80%, poor per-GPU performance	MFU (Model FLOPs Utilization)	Kernel-level levers (section A.5), mixed precision, program goodput levers (section 9.7)
Communication (C_2)	Network bandwidth ($\text{BW}_{\text{network}}$)	High AllReduce time, low scaling efficiency, communication > 30% of step	Scaling efficiency (η_{scaling}), communication fraction ($T_{\text{comm}}(N)/T_{\text{step}}(N)$)	Gradient compression, overlap compute/communication, topology optimization (chapter 5)
Coordination (C_3)	Synchronization overhead and failure recovery	Low runtime goodput, frequent restarts, large pipeline bubbles, scheduler churn	Runtime goodput ($\eta_{\text{runtime}} \approx T_{\text{useful}}/T_{\text{wall}}$)	Async checkpointing, elastic training, faster failure detection (chapter 7)

B.3 The Fleet Law

Under ideal partitioning of fixed work, the same classification can be written as a time budget for each distributed step. The Fleet Law, introduced in section 1.5.4, decomposes every distributed training step into the distributed-step time budget:

$$T_{\text{step}}(N) = \frac{T_{\text{compute}}}{N} + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}}$$

This equation is the fleet-scale counterpart of the iron law. Where the iron law decomposes single-machine execution into data movement, compute, and overhead, the Fleet Law decomposes distributed execution into local arithmetic, network data transfer, and synchronization logic. The diagnostic strategy is identical: measure each term, identify which dominates, and direct engineering effort at the dominant term.

B.3.1 Component decomposition

Each Fleet Law term maps to specific measurable activities that can be isolated in a timeline trace or cluster telemetry log. By attributing every millisecond of step duration to one of four terms, engineers can tell whether hardware is doing useful math, moving bytes across wires, or waiting on barriers:

- $T_{\textbf{compute}}/N$: Forward pass, backward pass, and optimizer step, all local arithmetic after distribution across N devices. Governed by MFU and per-GPU kernel efficiency. Improvements come from better kernels, mixed precision, and operator fusion.
- $T_{\textbf{comm}}(N)$: AllReduce of gradients, AllGather of parameters (in FSDP/ZeRO), activation transfers in tensor/pipeline parallelism. Governed by network bandwidth and collective algorithm choice. Improvements come from gradient compression, hierarchical collectives, and compute-communication overlap.
- $T_{\textbf{sync}}(N)$: Synchronization barriers, checkpoint writes, failure detection and recovery, pipeline bubble idle time, scheduler preemptions, and maintenance windows. Governed by cluster reliability and orchestration software. Improvements come from asynchronous checkpointing, elastic training, and faster failure detection.
- $T_{\textbf{overlap}}$: Communication or coordination time hidden behind useful arithmetic. Governed by scheduling and implementation overlap.

Operational events may be amortized into the step budget or accounted for separately in runtime goodput. The same checkpoint or recovery delay must not appear in both. The compute-time fraction follows:

$$f_{\text{compute}} = \frac{T_{\text{compute}}/N}{T_{\text{step}}(N)}$$

When f_{compute} drops below 0.5, the fleet spends more time on communication and coordination than on useful arithmetic. This compute-time fraction is not total fleet efficiency; useful fleet efficiency also depends on MFU, scaling efficiency, and goodput. The C^3 taxonomy identifies which noncompute term is responsible.

B.4 C^3 Intersections

In a production cluster, a stalled training step rarely comes from a single isolated resource; a slow collective can idle tensor cores while a synchronous checkpoint write blocks both the GPUs and the network fabric. The Venn diagram in figure B.1 maps how distributed execution spans the node substrate (C_1), the networking backbone (C_2), and the orchestration layer (C_3). Because real bottlenecks arise at the pairwise and central boundaries where two or more of these dimensions compound, diagnosing them requires mapping observed symptoms back to the governing intersection and the engineering techniques that address it (table B.3).

Table B.3: C^3 Intersection Reference: Each zone maps specific techniques to the axes they span and the chapters that cover them. The pairwise intersections require reasoning about two domains simultaneously; the center requires all three.

Zone	Name	Key techniques	Book coverage
Compute (C_1)	The Node Substrate	Operator fusion, mixed precision (FP8/BF16), MFU optimization, local NVMe staging	chapter 2, chapter 4, chapter 9, section A.5
Communication (C_2)	The Networking Backbone	Rail-optimized Clos topologies, RoCEv2/InfiniBand fabrics, collective algorithms	chapter 3, chapter 5
Coordination (C_3)	The Orchestration Layer	Gang scheduling (Slurm, Kubernetes, Ray), state consensus, heartbeat failure detection	chapter 7, chapter 8, chapter 12

Zone	Name	Key techniques	Book coverage
Compute ∩ Communication ($C_1 \cap C_2$)	Hiding Network Latency	1F1B pipelining, CUDA stream compute-comm overlap, comm-to-compute ratio (ρ) tuning, wire compression	chapter 5, chapter 6
Compute ∩ Coordination ($C_1 \cap C_3$)	Preserving Progress Against Failure	Non-blocking async checkpointing, pipeline bubble reduction, in-memory recompute vs. rollback, fast node drain	chapter 7, chapter 8
Communication ∩ Coordination ($C_2 \cap C_3$)	Collective Synchronization	Straggler mitigation, hierarchical collectives (NVLink vs. IB), bounded staleness, topology-aware rank mapping	chapter 5, chapter 7
Compute ∩ Communication ∩ Coordination ($C_1 \cap C_2 \cap C_3$)	Fleet ML Systems Engineering	Fleet Law decomposition, 4D parallelism (DP+TP+PP+EP), memory sharding (FSDP2, ZeRO-3), ML productivity goodput	chapter 6, chapter 10, chapter 11

B.4.1 Compute ∩ communication

When a GPU finishes its backward pass in 40 ms but waits 60 ms for an AllReduce collective to return gradients over the network, its tensor cores sit idle even though both the compute kernels and the links are working as specified. This intersection governs whether the system can hide communication behind computation. The communication-to-compute ratio ($\rho = T_{\text{comm}}(N)/(T_{\text{compute}}/N)$) is the key diagnostic metric (section C.3.3). When $\rho < 1$, computation takes longer than communication and there is enough arithmetic duration to hide the transfer in principle; dependencies and resource contention determine whether that overlap is achievable. When $\rho > 1$, GPUs complete their local work before the network delivers the next round of operands, and the system becomes communication-bound.

In a GPU trace, this bottleneck appears as exposed communication intervals where tensor cores sit idle awaiting collective completion (such as `ncclKernel_AllReduce` or `ncclKernel_AllGather`). Engineering at this intersection focuses on overlap and compression. Engineers launch AllReduce during the backward pass before all gradients are finalized, pipeline activation transfers in 1F1B schedules, and increase microbatch sizes or gradient accumulation steps to dilute communication frequency. They also apply wire quantization to reduce payload bytes, while accounting for metadata, codec time, and numerical error (chapter 5, chapter 6).

B.4.2 Compute ∩ coordination

A 1,024-GPU cluster that pauses every thirty minutes to flush a multi-terabyte checkpoint to shared storage turns expensive arithmetic capacity into idle wait time. This intersection captures the idle compute caused by orchestration and resilience mechanisms. While compute engines should remain saturated with arithmetic, coordination overhead repeatedly forces hardware into idle states. Three distinct mechanisms govern this intersection:

1. **Pipeline bubbles**: In pipeline parallelism (PP), stages idle during pipeline fill (warmup) and drain (cooldown). For N_{PP} pipeline stages and m microbatches, the bubble fraction $b_{\text{pipe}} = (N_{\text{PP}} - 1)/(N_{\text{PP}} - 1 + m)$ (section D.4.1) is accelerator time lost to idling; more microbatches or interleaved schedules shrink it.
2. **Checkpoint stalls**: Synchronous checkpoint writes that block the training loop to serialize terabytes of weights and optimizer states convert expensive compute time into storage wait time. Asynchronous checkpointing mitigates this by snapshotting host-pinned or local Non-Volatile Memory Express (NVMe) memory in milliseconds, delegating remote persistence to background I/O threads.
3. **Recovery rollback**: When a node fails in an N-device cluster, training halts and rolls back to the last valid checkpoint, squandering the arithmetic executed since that snapshot. Minimizing the failure blast radius through fast health heartbeats and dynamic node replacement bounds this compute loss (chapter 7, chapter 8).

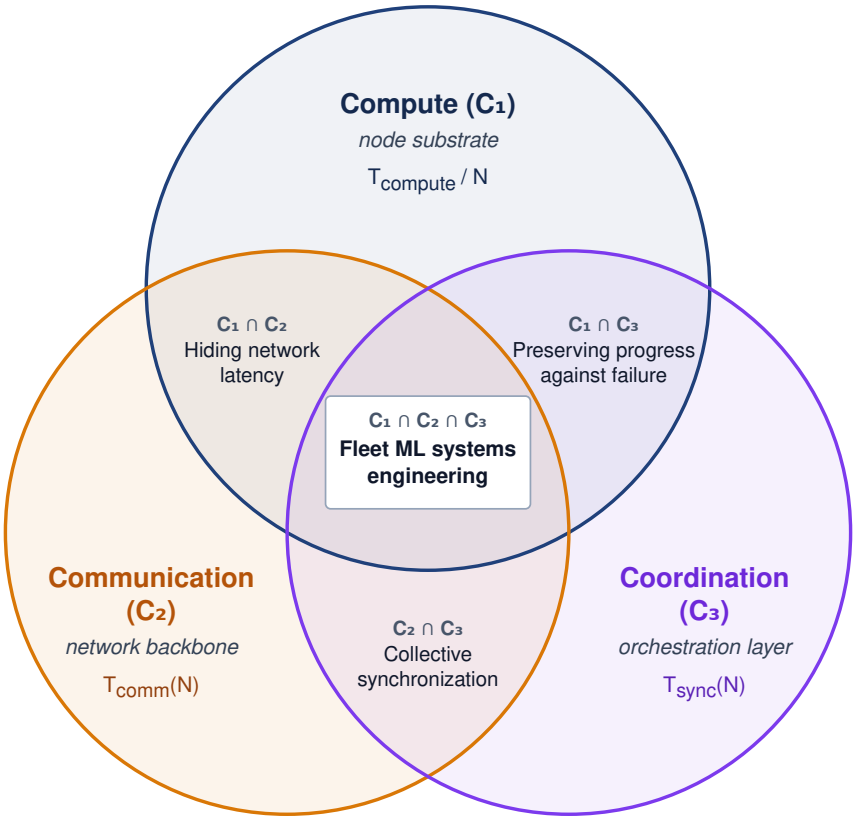

Figure B.1: The C^3 Intersection Landscape: Each circle is one operational dimension (Compute, C_1, the node substrate; Communication, C_2, the network backbone; Coordination, C_3, the orchestration layer), labeled with the Fleet Law term it contributes to step time. Each pairwise overlap names the interaction between two dimensions (hiding network latency behind computation, preserving training progress against failure, and synchronizing collectives). At the center, where all three meet, is fleet ML systems engineering.

B.4.3 Communication ∩ coordination

When a single GPU or network link drops packets or throttles its clock during an AllReduce collective, every other rank in the group must wait at the dependency boundary before the next step can proceed. This intersection captures the synchronization cost embedded in network communication. A collective operation (AllReduce, AllGather, ReduceScatter) moves data over physical wires (C_2) while creating dependencies among its participating ranks (C_3). It need not impose a separate cluster-wide barrier, but a delayed contribution can prevent dependent outputs from becoming ready.

The governing failure mode is that one slow participant paces the whole collective. If a single GPU or link experiences packet drops, thermal throttling, or ECC retry delays, every other GPU in the collective waits, transforming a localized degradation into a cluster-wide halt. At fleet scale ($N \geq 10{,}000$), tail latency dominates the barrier time, because the probability that at least one node straggles approaches unity ($1-(1-p_{\text{straggle}})^N \to 1$). Engineering at this intersection focuses on decoupling synchronization from raw communication. Hierarchical collectives isolate high-bandwidth intra-node NVLink domains from inter-node InfiniBand fabrics. Bounded-staleness asynchronous training absorbs transient tail latency. Topology-aware gang scheduling aligns collective communication rings with physical spine-leaf switch rails (chapter 5, chapter 7).

B.4.4 Fleet-scale ML systems engineering

Partitioning a 405-billion-parameter model across 16,384 GPUs requires fitting shard states into high-bandwidth memory (HBM) (C_1), matching tensor and pipeline traffic to intra-node NVLink and inter-node InfiniBand bandwidths (C_2), and keeping pipeline bubbles and checkpoint stalls from consuming wall-clock time (C_3). At the center ($C_1 \cap C_2 \cap C_3$), all three constraints converge simultaneously. This is the discipline of **fleet-scale ML systems engineering**. Systems engineers at this center do not optimize a single accelerator or a single network switch; they manage the trade-offs of the distributed machine as an integrated whole. The central engineering challenges encompass three coupled design problems:

- **4D parallelism co-design**: Balancing data parallelism (DP), tensor parallelism (TP), pipeline parallelism (PP), and expert parallelism (EP) so that memory consumption fits within accelerator HBM (C_1), high-frequency communications map to intra-node NVLink and low-frequency communications map to inter-node fabrics (C_2), and pipeline bubbles and barrier dependencies are minimized (C_3).
- **Memory sharding vs. communication volume**: Frameworks like PyTorch FSDP2 and DeepSpeed ZeRO-3 eliminate memory redundancy (C_1) by partitioning model states across workers, but pay for it with repeated AllGather communication (C_2) and state coordination metadata (C_3).
- **ML productivity goodput**: ML productivity goodput (MPG), the fraction of a fleet's capacity-time that becomes useful work, is the product of scheduling goodput ($\eta_{\text{scheduling}}$), runtime goodput (η_{runtime}), and program goodput (η_{program}) (section 9.2.4). These factors cut across the C^3 axes rather than mapping one to one onto them. Communication overlap and shorter pipeline bubbles improve program goodput, while checkpointing and recovery, which are Coordination costs, determine runtime goodput.

B.4.5 C^3 coordination and overlap

Just as single-node execution transforms the additive iron law into a maximum bounded by dependencies (section A.2.1), overlap at fleet scale transforms the Fleet Law (equation 1.3) by maximizing T_{overlap}. Pipeline schedules, separate communication streams, asynchronous checkpoints, hierarchical collectives, and double-buffered data streaming each hide one C^3 axis behind another (table B.4).

Table B.4: C^3 Overlap Techniques: System techniques that hide communication or coordination behind useful computation, maximizing T_{overlap} in the Fleet Law.

Technique	C^3 overlap	Mechanism	Implementation
1F1B Pipeline Schedule	C_1 overlaps C_2	Interleaves forward and backward microbatches to overlap cross-stage activation transfer with arithmetic	Microbatch queue scheduling (chapter 6)
CUDA Stream Decoupling	C_1 overlaps C_2	Launches non-blocking AllGather or AllReduce on separate communication streams concurrently with GEMMs	Dedicated NCCL streams (chapter 5)
Async Checkpointing	C_1 overlaps C_3	Fast in-memory or local NVMe snapshot while GPU arithmetic resumes immediately	Background host thread flushes to storage (chapter 4, chapter 7)
Hierarchical Collectives	C_2 overlaps C_3	Partitions collective into local NVLink reduction and inter-node fabric transfer, isolating tail jitter	Hierarchical NCCL communicators (chapter 5)
Double-Buffered Data Streaming	C_1 overlaps C_2	Prefetches and decompresses next batch shard in pinned host memory while GPUs compute current batch	Host pinned memory + async H2D transfer (chapter 4)

B.4.6 Storage as the persistent substrate

A common diagnostic question is why Storage does not appear as a fourth explicit coordinate in the C^3 taxonomy or as an independent T_{storage} term in the Fleet Law. The physical data center is procured

around Compute, Network, and Storage (chapter 4), yet at runtime storage devices execute no model arithmetic. Storage overhead therefore manifests dynamically through the three C^3 dimensions:

- **In Compute (C_1)**: High-throughput local NVMe drives function as the tier-3 extension of the accelerator memory hierarchy (below SRAM and HBM), serving parameter offloading (ZeRO-Offload) and activation spilling.
- **In Communication (C_2)**: Remote dataset streaming from parallel file systems (Lustre, GPFS) or object stores (S3, Ceph) shares cluster network bandwidth and network interface card (NIC) capacity with collective communication traffic.
- **In Coordination (C_3)**: Durable distributed storage provides the persistent state anchor required for checkpoint serialization, node failover, and job rescheduling.

When storage latency stalls a training run, profilers identify the symptom not as raw disk latency, but either as a network communication bottleneck (C_2, dataset starvation) or as a blocking coordination barrier (C_3, synchronous checkpoint flush). Storage is the persistent substrate supporting all three dimensions, rather than a separate term in the step execution budget.

B.5 Rules of Thumb

In the middle of a production incident, fast heuristics narrow the search space before a profiler is needed. These thresholds provide that first line of defense.

B.5.1 The C^3 traffic light

Each C^3 axis has one headline metric, and its value against two thresholds sorts the axis into green, yellow, or red (table B.5). Comparing observed MFU, communication fraction, and runtime goodput against these bands immediately isolates which subsystem warrants a deeper trace.

Table B.5: C^3 Traffic Light: Quick triage thresholds for fleet-scale diagnosis. Green means the axis is healthy; yellow means it deserves investigation; red means it is the likely bottleneck. These thresholds assume well-optimized large-model training on current-generation hardware.

C^3 axis	Green (healthy)	Yellow (investigate)	Red (bottleneck)
Compute	MFU $>$ 50%	MFU 30%–50%	MFU $<$ 30%
Communication	Comm fraction $<$ 20%	Comm fraction 20–40%	Comm fraction $>$ 40%
Coordination	Runtime goodput $>$ 90%	Runtime goodput 75–90%	Runtime goodput $<$ 75%

B.5.2 The bottleneck diagnostic table

An optimization pays only when it shortens the term that binds: upgrading to faster GPUs in a communication-bound fleet or adding network bandwidth to a compute-bound fleet leaves step time flat (table B.6).

Table B.6: What Works vs. What Is Wasted at Fleet Scale: Optimizing a nondominant C^3 axis yields limited improvement and often worsens cost efficiency. A communication-bound fleet will gain little from faster GPUs until AllReduce and other communication bottlenecks are addressed.

If the fleet is…	Dominant term	Optimization that works	Optimization that is wasted
Compute-bound	T_{compute}/N	Better kernels, mixed precision, operator fusion, next-gen accelerators	More network bandwidth (GPUs are not waiting on the network)
Communication-bound	$T_{\text{comm}}(N)$	Gradient compression, compute-comm overlap, hierarchical collectives, InfiniBand upgrade	Faster GPUs (they spend a larger share of each step waiting for the network)
Coordination-bound	$T_{\text{sync}}(N)$	Async checkpointing, elastic training, faster failure detection, fewer pipeline stages	Faster arithmetic alone; faster networks help only transfer-related coordination costs

B.6 C^3 Case Studies

In production, two clusters with the same slow step time can require opposite fixes depending on whether the stall comes from low per-GPU kernel utilization, exposed gradient synchronization, or frequent recovery restarts. The three cases below isolate one dominant C^3 axis at a time and trace the diagnosis from initial symptom to targeted fix.

B.6.1 Case 1 on the underutilized compute fleet

Symptom

An engineering team provisions 4,096 H100 GPUs for a large language model training run. The training loop runs without errors, but the PyTorch Profiler shows MFU of only 15 percent. The network profiler shows communication accounts for less than 10 percent of step time. The cluster is running but barely working.

Diagnosis

The **Compute** axis is the leading hypothesis to investigate. With 15 percent MFU, 85 percent of the fleet's peak capacity is not credited to useful model FLOPs. Low exposed communication makes local execution a useful starting point; kernel traces and host-side waits must distinguish inefficient arithmetic from data starvation or unobserved coordination gaps.

The fix

If traces confirm inefficient local execution, target the compute axis across 4,096 accelerators to raise per-device kernel utilization. Three adjustments recover most of the missing arithmetic throughput:

- **Mixed precision**: Ensure BF16/FP8 Tensor Cores are engaged. A common culprit is FP32 fallback in normalization layers or loss computation.
- **Operator fusion**: Use `torch.compile` or similar just-in-time (JIT) compilation to fuse element-wise operations and reduce kernel launch overhead.
- **Batch size tuning**: If per-GPU batch size is too small, the matrix multiplications have insufficient arithmetic intensity to saturate the Tensor Cores.

Raising MFU from 15 percent to 50 percent on the same hardware delivers 50 percent/15 percent = 3.3× more useful work, the equivalent of tripling the fleet without buying a single GPU. Because communication is already below 10 percent of step time, all of that per-device gain translates directly into faster step completion.

B.6.2 Case 2 on the communication wall

Symptom

A 512-GPU data-parallel training run achieves 45 percent illustrative compute-phase efficiency during compute kernels on each GPU, which is good per-device kernel efficiency. Wall-clock MFU over the full step (including AllReduce) is only 20.2 percent because AllReduce consumes 55 percent of every training step. Scaling from 64 GPUs to 512 GPUs yields only 4× speedup instead of the expected 8×, or 50 percent scaling efficiency.

Diagnosis

The **Communication** axis dominates. Each GPU has healthy compute-phase efficiency, but more than half the step time is spent synchronizing gradients across the network. The system is **Communication-Bound**: adding more GPUs can worsen the exposed communication fraction. Under fixed-work scaling, local computation shrinks while the collective retains startup and transfer costs.

The fix

Target the communication axis without touching the per-GPU computation. Three network-level interventions shrink or hide the exposed collective time:

- **Compute-communication overlap**: Launch AllReduce during the backward pass rather than waiting until it completes. Modern frameworks (FSDP, DeepSpeed) support this natively.

- **Gradient compression**: Apply Top-k sparsification or quantization to reduce the bytes crossing the network by 10–100×.
- **Hierarchical collectives**: Use intra-node NVLink for the first reduction stage, then inter-node InfiniBand only for cross-node aggregation, reducing inter-node participants relative to a flat rank-level algorithm. Actual traffic savings depend on the collective layout, NIC allocation, and topology.

If communication were eliminated entirely, throughput would increase by 2.2× (Amdahl's Law applied to the 45 percent compute fraction). Realistically, reducing communication from 55 percent to 20 percent of step time would recover most of the lost scaling.

B.6.3 Case 3 on coordination overhead

Symptom

A 10,000-GPU training run shows 40 percent compute-phase model FLOPs utilization per device, measured only while kernels execute, and communication accounts for only 15 percent of step time. Both figures look healthy. Yet only 60 percent of wall time is left after the run's coordination overheads. That share is not a goodput factor, because it mixes losses that belong to different factors. Pipeline bubbles lie inside retained steps and are a program goodput loss, while checkpoints, recovery, maintenance, and preemption fall outside the steps (section 9.2.3). The remaining 40 percent of wall time is consumed by checkpoint writes, failure recovery restarts, pipeline bubble idle time, scheduler preemptions (17 percent of wall time), and maintenance windows.

Diagnosis

The **Coordination** axis dominates. Although per-GPU computation and inter-node communication are both efficient, nonproductive overhead consumes 40 percent of wall time. Scheduler preemptions take 17 percent, and failure recovery takes 10 percent (because at 10,000 GPUs, GPU failures occur about every 5 hours). Pipeline bubbles take 5 percent, checkpoint writes take 3 percent, and maintenance windows take 5 percent. Because faster arithmetic alone cannot recover this lost time, checkpoint bandwidth, recovery mechanisms, and scheduling must be evaluated against the specific overhead each can reduce.

The fix

Target the coordination axis so that orchestration and fault-recovery events stop idling the healthy accelerators. Four operational changes attack both the in-step bubble losses and the out-of-step downtime:

- **Asynchronous checkpointing**: Overlap checkpoint writes with the next training step, reducing visible checkpoint overhead from 3 percent to near zero.
- **Elastic training**: When a node fails, shrink the job and continue rather than halting all 10,000 GPUs for recovery. This converts the 10 percent failure recovery cost into a smaller throughput reduction.
- **Pipeline schedule optimization**: Switch from GPipe to an interleaved 1F1B schedule to reduce bubble fraction, or increase microbatch count per pipeline flush.
- **Faster failure detection**: Reduce heartbeat timeout from 30 seconds to 5 seconds with hardware-level health monitoring, cutting the idle time between failure occurrence and recovery initiation.

B.7 Production Troubleshooting

Each common fleet-scale failure mode first shows up as a user-visible symptom, and that symptom points to the C^3 axis most likely responsible (sometimes two), a measurement that confirms it, and an action (table B.7). Operators can match an observed cluster symptom to its primary C^3 axis, run the listed diagnostic check, and apply the corresponding engineering fix.

Table B.7: C^3 Troubleshooting Matrix: Root cause identification and remediation for common fleet-scale bottlenecks. Each row connects a user-visible symptom to the C^3 axis most likely responsible, reducing the search space before reaching for a profiler.

Symptom	C^3 axis	Diagnostic question	Measurement	Action
Low MFU despite fast network	**Compute**	Are Tensor Cores engaged? Is batch size sufficient for arithmetic intensity?	Per-GPU kernel trace (Nsight/PyTorch Profiler)	Enable mixed precision, increase per-GPU batch size
Throughput plateaus when adding GPUs	**Communication**	Does AllReduce time grow faster than computation shrinks?	NCCL trace, ρ ratio	Gradient compression, hierarchical collectives, overlap
Frequent job restarts	**Coordination**	What is the cluster MTBF? Is detection fast enough?	Failure logs, MTBF calculation	Elastic training, faster detection, smaller blast radius
High GPU-hours but slow progress	**Coordination**	What fraction of GPU-hours produce useful training steps?	Runtime goodput ($T_{\text{useful}}/T_{\text{wall}}$)	Async checkpointing, reduce pipeline stages, eliminate scheduler churn
Scaling efficiency drops with cluster size	**Comm/Coord**	Is the bottleneck network bandwidth or synchronization barriers?	Separate $T_{\text{comm}}(N)$ from $T_{\text{sync}}(N)$	If comm: compress or overlap. If coord: async methods
Stragglers slow entire job	**Comm ∩ Coord**	Is one node consistently last to reach the AllReduce barrier?	Per-node step time histogram	Straggler detection + replacement, bounded staleness, backup workers

B.8 Tooling Map

Confirming a candidate C^3 bottleneck requires inspecting the layer where that constraint lives: per-GPU kernel profilers for Compute, collective and fabric counters for Communication, and scheduler and goodput logs for Coordination (table B.8).

Table B.8: C^3 Tooling Map: Profiling utilities for diagnosing fleet-scale bottlenecks. Start with the primary tool for quick triage; use secondary tools for deep-dive analysis. Compute tools operate per-GPU; Communication tools operate at the network layer; Coordination tools operate at the cluster/job level.

C^3 axis	Key metric	Primary tool	Secondary tool
Compute	MFU, kernel utilization	PyTorch Profiler (TensorBoard plugin)	Nsight Compute (per-kernel roofline analysis)
Communication	AllReduce time, ρ ratio	NCCL debug logs (`NCCL_DEBUG=INFO`)	Nsight Systems (timeline), `ibstat/perfquery` (IB)
Coordination	Runtime goodput, restart count	Cluster scheduler logs (Slurm, K8s event logs)	Custom goodput dashboards (for example, Google ML Goodput)

B.9 C^3 Scorecard

Once profiling tools supply the raw counters, three fleet-level efficiency ratios extend the single-machine scorecard (section A.9) to grade distributed cluster health against production thresholds (table B.9).

Table B.9: The C^3 Efficiency Rubric: Use these three numbers to characterize fleet health. The MFU denominator uses aggregate fleet peak, written as NR_{peak} when R_{peak} is a per-device value. These illustrative thresholds guide investigation; passing all three does not prove that inexpensive optimizations are exhausted.

C^3 axis	Metric	Definition	Failing grade	Passing grade
Compute	**MFU**	$\frac{O_{\text{useful}}}{NR_{\text{peak}} \times T_{\text{step}}}$	$< 30\%$	$> 50\%$
Communication	**Scaling Efficiency** (η_{scaling})	$\frac{T_1}{N \times T_N}$	$< 56\%$	$> 87\%$

C^3 axis	Metric	Definition	Failing grade	Passing grade
Coordination	**Runtime Goodput** (η_{runtime})	$\frac{T_{\text{useful}}}{T_{\text{wall}}}$, retained steps at the job's fault-free step rate over held time	$< 75\%$	$> 90\%$

B.10 C^3 Infrastructure Diagnostic Decision Tree and Fleet Diagnostic Scorecard

When step time suddenly doubles in the middle of a training run, on-call operators see raw hardware counters—elevated GPU temperatures, PCIe retry alerts, or InfiniBand pause frames—rather than a neatly labeled C^3 axis. Cluster operations teams need a systematic procedure for mapping those low-level hardware signals to the binding C^3 axis before taking remedial action.

B.10.1 Telemetry signal mapping to C^3 axes

A degraded PCIe riser or a clogged cooling loop first surfaces in hardware counters across silicon, PCIe buses, thermal sensors, and network switches. Monitoring daemons collect four primary classes of low-level telemetry to catch these faults early:

- **MFU Drops**: Sudden or sustained drops in model FLOPs utilization below expected baselines.
- **PCIe AER (Advanced Error Reporting) Retries**: Non-fatal bus errors causing packet retries, link downgrades (e.g., PCIe Gen 5 $\rightarrow$ Gen 3), or host-to-device direct memory access (DMA) throttling.
- **Thermal Throttling**: GPU SM clock downclocking (e.g., 1.9 GHz $\rightarrow$ 1.1 GHz) triggered by junction thermal limits or cooling fan degradation.
- **InfiniBand Credit Stalls and PFC Pause Frames**: Priority flow control (PFC) pause frames, congestion notification packets (CNP), or switch buffer credit starvation.

B.10.2 Diagnostic decision tree

Triage starts from whichever telemetry source shows the anomaly and narrows to the specific signal before naming the binding axis (figure B.2). A node-side signal points to the compute axis, while a network-side signal separates fabric congestion, a communication cost, from a single straggler that holds every rank at the barrier, a coordination cost.

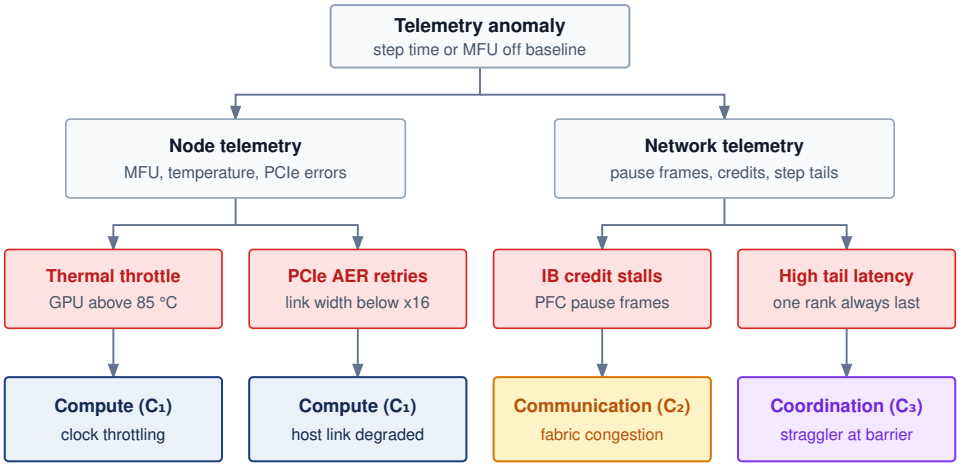

Figure B.2: Telemetry-First Triage: An anomaly is traced first to the telemetry source that shows it, then to the specific signal, and only then to the C^3 axis that binds. Thermal throttling and PCIe link degradation both lower per-device throughput and so point to compute, credit stalls point to communication, and a single rank that is always last points to coordination. The 85 °C threshold is illustrative.

B.10.3 Structured diagnostic telemetry table

Once the decision tree isolates the anomalous hardware or OS counter, pairing that counter with its physical subsystem lets automated health checks route the alert directly to the appropriate repair workflow (table B.10).

Table B.10: C^3 Telemetry Diagnostic Mapping: Direct mapping from silicon-level telemetry signals to the binding C^3 axis and its remediation.

Telemetry signal	Hardware/OS counter	Root cause	C^3 axis	Remediation action
MFU Drop with High GPU Temp	`gpu_temperature > 85°C`, `clocks_throttle_reasons.sw_thermal`	Thermal throttling reduces GPU SM clock speed	**Compute**	Check coolant flow, replace thermal interface material (TIM), adjust fan curves
MFU Drop with Low Host-GPU BW	`pcie_aer_correctable_errors`, `pcie_link_width < x16`	PCIe link degradation/retries due to signal noise	**Compute**	Reseat GPU/PCIe riser card, clean PCIe edge connectors, replace cable
AllReduce Stalls with Pause Frames	`PFC_PAUSE_XOFF_FRAMES`, `infiniband_credit_stalls`	Network fabric congestion, IB buffer exhaustion	**Communication**	Enable adaptive routing (AR), rebalance ECMP hashing, tune NCCL buffer size
NCCL Barrier Timeout/Hang	`NCCL_DEBUG=INFO` log `Unhandled CUDA error`, `heartbeat_miss`	Single-node straggler or silent kernel hang	**Coordination**	Drain degraded node, isolate via node health check script, trigger fast restart
High Step Overhead with Idle GPU	`nvlink_throughput == 0`, `fs_write_latency > 1s`	Blocking synchronous checkpoint write to storage	**Coordination**	Switch to asynchronous double-buffered checkpointing or host RAM staging

B.10.4 Fleet diagnostic scorecard

The fleet diagnostic scorecard in table B.11 carries each single-node Data, Algorithm, and Machine (D·A·M) concern to the fleet telemetry that exposes it, the C^3 axes it affects, and a target benchmark. Each row links a single-node resource category to the cluster-level telemetry signal and target benchmark that govern distributed execution.

Table B.11: Fleet Diagnostic Scorecard: Each single-node D·A·M concern maps to the fleet telemetry that exposes it and to the C^3 axes it affects.

D·A·M axis	Fleet telemetry signal	C^3 axis	Target efficiency benchmark
Data (D)	Network fabric RX/TX saturation, storage write I/O wait times	**Communication/Coordination**	I/O wait time $< 2\%$ of step time; NCCL ring bandwidth $> 85\%$ of fabric peak
Algorithm (A)	Pipeline bubble idle fraction, partitioning imbalance across stages	**Compute/Coordination**	Pipeline bubble fraction $< 20\%$ under the balanced-stage model
Machine (M)	Model FLOPs utilization (MFU), PCIe AER errors, thermal and power throttling	**Compute**	MFU $> 50\%$; thermal throttle events $= 0$; PCIe link width held at maximum (Gen 5 x16)

B.11 Scaling Laws Through the C^3 Lens

Scaling laws are usually written in algorithmic FLOPs, but a training fleet delivers only the useful FLOPs that survive utilization, communication, and coordination losses. Translating a scaling-law target into C^3 terms means removing its hidden perfect-systems assumption and replacing peak FLOP/s with effective FLOP/s, the usable throughput that connects model-quality forecasts to real cluster capacity.

B.11.1 The ideal-system assumption in scaling laws

Scaling laws (Kaplan et al. 2020; Hoffmann et al. 2022) and their successors predict model quality as a function of algorithmic training compute. They usually abstract away systems efficiency: wall-clock time, accelerator peak FLOP/s, compute-phase efficiency, communication overhead, and scheduler

losses appear later when teams provision hardware to deliver the target compute budget. In C^3 terms, scaling-law FLOPs must be converted into raw fleet capacity after compute-phase, distributed-step, and runtime-goodput losses.

The gap between scaling-law predictions and observed training outcomes is, in large part, a C^3 gap. A team that budgets 10^{24} peak hardware FLOPs of capacity will deliver fewer useful model FLOPs after three multiplicative losses: per-GPU compute-phase efficiency, inter-node scaling efficiency (η_{scaling}), and runtime goodput.

B.11.2 The effective FLOP/s concept

A fleet's **effective FLOP/s** is the usable throughput after compounding three disjoint C^3 efficiency factors across the cluster:

$$R_{\text{eff}} = (N R_{\text{peak}}) \times \underbrace{\text{MFU}}_{\text{Compute}} \times \underbrace{\eta_{\text{scaling}}}_{\text{Distributed step}} \times \underbrace{\eta_{\text{runtime}}}_{\text{Runtime goodput}}$$

MFU, the single-device baseline, measures useful model arithmetic against peak throughput only while local computation is active. This compute-phase factor is not the step-averaged fleet MFU of table B.9, whose T_{step} denominator already includes the in-step losses that η_{scaling} counts, so the product does not count them twice. Scaling efficiency captures communication and other distributed-step losses, including pipeline bubbles. Runtime goodput captures time outside those measured steps, including checkpointing, failure recovery, scheduler preemptions, and maintenance. These disjoint measurement windows prevent the same loss from being counted twice, while a concrete estimate makes the multiplicative loss visible at fleet scale.

> **Systems Perspective B.1: Compound C^3 loss on a 100,000-GPU cluster**
>
> Consider a 100,000-GPU H100 cluster with 98,900 PFLOP/s of peak aggregate throughput. After the three disjoint C^3 efficiency factors:
>
> $$\text{Effective} = 98{,}900 PFLOP/s \times 0.50 \times 0.56 \times 0.65 \approx 17999.8 PFLOP/s$$
>
> The fleet delivers 18.2 percent of its peak capacity as useful training work. The **compound C^3 loss**, the ratio of peak to effective compute, is 5.5×, so achieving a given effective compute budget requires 5.5× the raw hardware. The compute phase retains a 50 percent efficiency factor, and the distributed step retains a 56 percent scaling factor (using the modeled 8,192-GPU tier as an illustrative proxy that includes pipeline losses). Runtime goodput retains a 65 percent factor after checkpoints, failures, scheduler preemptions, and maintenance.
>
> These losses reflect the physics of fleet-scale computation rather than a failure of engineering. The C^3 taxonomy quantifies where the losses occur so that optimization effort targets the dominant term.

B.12 Summary

The C^3 taxonomy provides a systematic framework for diagnosing fleet-scale bottlenecks. Each axis maps to a distinct physical constraint: arithmetic throughput and MFU bound compute; network bandwidth and collective algorithm efficiency bound communication; and synchronization overhead, failure recovery, and operational losses bound coordination. The Fleet Law quantifies these constraints, enabling systematic diagnosis. Use the C^3 traffic light for quick triage, the bottleneck diagnostic table to choose the right lever, and the C^3 scorecard to grade fleet health.

> **Key Takeaways: Where to look first at fleet scale**
>
> - **Every fleet-scale bottleneck has a dominant C^3 axis**: Compute, Communication, or Coordination, with many real bottlenecks spanning intersections. Identify the dominant axis before optimizing.

- **Measure the C^3 scorecard**: Compare MFU, scaling efficiency, and goodput with workload-specific baselines. The displayed thresholds screen for problems; traces must confirm the cause before selecting an optimization.
- **The compound C^3 loss is multiplicative**: Peak FLOP/s $\times$ compute-phase efficiency $\times$ scaling efficiency $\times$ runtime goodput = effective FLOP/s. The illustrative 100,000-GPU case retains ~18.2 percent of peak.
- **Coordination is the new axis**: Single-node overhead remains relevant, while distribution adds cross-machine dependencies. The coordination case attributes 40 percent of its modeled budget to checkpoints, failures, pipeline bubbles, scheduling, and maintenance.
- **Optimizing the wrong C^3 axis yields limited improvement**: Faster GPUs cannot fix a communication-bound fleet by themselves; faster networks help coordination only where transfer time, such as checkpoint I/O, contributes to that overhead.

C

Fleet Foundations

Single-node performance engineering centers on memory hierarchy latencies and roofline ridge points. The moment a job spans multiple machines, network latency replaces cache latency, component failure becomes a statistical certainty, and interconnect bandwidth dictates scaling efficiency. This appendix gathers the fleet-scale constants, scaling laws, and thermal limits that underpin the book's multi-node chapters, giving readers the physical baselines needed to reason about cluster-scale infrastructure. In C^3 terms, these numbers define the physical scale on which compute capacity, communication bandwidth, and coordination overhead become measurable.

How to Use This Appendix

This appendix is designed as a reference when translating fleet-scale design questions into specific physical constraints and identifying available optimization levers. The five routing bullets below map common engineering questions to the sections that follow.

- **When evaluating inter-node communication**: Consult section C.2 for the bandwidth and latency hierarchy.
- **When sizing GPU allocations**: Consult section C.3 to analyze why doubling accelerator counts does not halve training time.
- **When calculating cluster reliability**: Check section C.2 for mean time between failures (MTBF) tables and failure probability models.
- **When evaluating power and cooling limits**: Consult section C.4 for power density ceilings and cooling constraints.
- **When budgeting runtime overhead**: Use section C.2 for the four overhead categories that erode useful wall time.

C.1 The Three System Paradigms

Machine learning infrastructure at scale inherits design DNA from two distinct computing lineages and then breaks the assumptions of both. A designer must understand where ML systems borrow from high-performance computing (HPC) and where they borrow from warehouse-scale computing (WSC) to choose the right design trade-offs. A system architect who treats ML training as "just HPC" will build infrastructure that cannot tolerate failures. One who treats it as "just web services" will build infrastructure that cannot sustain the tight coupling that synchronous training demands.

High-performance computing

HPC systems descend from the supercomputer tradition, where the design goal is to maximize FLOP/s on tightly coupled simulations such as weather modeling, molecular dynamics, and nuclear physics. Because every node participates in lockstep, these clusters rely on specialized interconnects (InfiniBand), low-latency fabrics, and homogeneous hardware. Their fault tolerance follows the **checkpoint/restart** model: if one node fails, the entire job stops, rolls back to the last checkpoint, and resumes from that saved state. Batch schedulers such as Slurm allocate rigid resource shapes ("512 nodes for 24 hours"), treating nodes as individually tracked "pets."

Warehouse-scale computing

WSC systems descend from the web services tradition, aiming instead to maximize queries per second across loosely coupled services such as search, email, and social media. Because requests are independent, these fleets mix hardware generations over commodity Ethernet and rely on **redundancy** for fault tolerance: when a node fails, a load balancer reroutes traffic to another replica without the user noticing. Dynamic schedulers such as Kubernetes and Borg bin-pack workloads elastically across interchangeable, expendable "cattle" nodes.

The ML fleet as hybrid architecture

ML systems require the computational throughput of HPC to train large models with synchronous gradient updates, yet they must operate at the scale and failure rates of WSC, where thousands of accelerators run for weeks. This tension produces a hybrid architecture that borrows selectively from both lineages.

Training workloads are synchronous and bandwidth-hungry like HPC jobs, whereas long-running fleet operations demand WSC-style fault tolerance, and inference workloads reverse the blend by pairing WSC latency sensitivity with HPC compute density. To serve both regimes, the network fuses TCP/IP control planes with InfiniBand or NVLink data planes. Fault tolerance combines checkpoint/restart with **elastic training** strategies where supported, which allow jobs to adjust resource allocations even when resharding, state restoration, or worker restarts remain necessary. Scheduling pairs all-or-nothing gang allocation for training with dynamic preemption and replacement. ML fleets therefore cannot adopt either HPC or WSC patterns wholesale; they combine elements from each according to the workload phase, as the design trade-offs across the three paradigms in table C.1 show.

Table C.1: Three System Paradigms: ML fleets inherit design DNA from both HPC and WSC but break assumptions of each. Training resembles HPC (tight coupling), inference resembles WSC (elastic serving), and fault tolerance is a hybrid of both.

Dimension	HPC (supercomputer)	WSC (web cloud)	ML fleet (AI cluster)
Philosophy	Maximize FLOP/s	Maximize QPS	Maximize Model Quality per Dollar/Watt
Coupling	Tight (MPI)	Loose (RPC/HTTP)	Hybrid (NCCL + RPC)
State	Stateful (RAM)	Stateless (DB-backed)	Semi-Stateful (Checkpoints + KV Cache)
Network	Latency-optimized	Bandwidth-optimized	Bisection-bandwidth critical
Bottleneck	Compute throughput (FLOP/s)	I/O (Disk/Net)	Fabric or HBM, by workload phase
Fault model	Checkpoint/Restart	Redundancy/Replicas	Checkpoint/restart + optional elasticity
Scheduling	Batch (Slurm)	Orchestration (K8s)	Gang + preemption
Node model	Pets (tracked)	Cattle (expendable)	Pets during job, cattle between jobs

Foundations recap

Five single-machine models reappear throughout the distributed chapters, each extended by a fleet-scale term. Appendix A states the single-machine forms; the list below records only what changes at scale.

- **Iron law** ($T \approx D_{\text{vol}}/\text{BW} + O/(R_{\text{peak}} \cdot \eta_{\text{hw}}) + L_{\text{lat}}$, section A.2): At fleet scale, the data-movement term expands to include inter-node communication, and the Fleet Law replaces the single-machine sum.
- **Roofline Model** (section A.3.1): At fleet scale, a third ceiling appears. **Network-bound** workloads are limited by inter-node bandwidth, which the communication-computation ratio in section C.3.3 measures.
- **Amdahl's Law**: Caps strong-scaling speedup at $1/r$ for serial fraction r (Amdahl 1967). At fleet scale, the serial fraction includes synchronization barriers, exposed collective communication, and pipeline bubbles (section C.3.2).

- **Training state rule** (section A.6): Mixed-precision training with the Adaptive Moment Estimation (Adam) optimizer holds about 16 bytes per parameter during training and about 14 bytes per parameter in a resumable checkpoint, before activations and framework metadata. At fleet scale, this footprint determines how model state is partitioned across devices (ZeRO, tensor parallelism, pipeline parallelism).
- **Little's Law** ($Q_{\text{req}} = \lambda_{\text{arr}} T_{\text{lat}}$) (Little 1961): Relates in-flight requests to arrival rate and time in system. At fleet scale, it sizes the number of serving replicas behind a load balancer (section F.1.1).

The transition from single-machine to fleet-scale reasoning requires extending these models with new dimensions: network topology, failure probability, and coordination overhead. The numbers in the next section provide the quantitative foundation for that extension.

C.2 Numbers Every Fleet Engineer Should Know

Crossing a server chassis boundary drops per-direction interconnect bandwidth by nearly an order of magnitude from intra-node NVLink to inter-node InfiniBand, while scaling a job from one node to thousands of accelerators shrinks the expected interval between hardware failures from years to hours. Just as single-machine analysis depends on a set of core numbers within each node, fleet-scale engineering is governed by the ratios and scaling behaviors that operate between nodes. Because hardware ratios shift across generations and failure rates depend on the assumed fault model, the reference values below provide baseline sanity checks while the accompanying equations show how to substitute different system assumptions.

> **Systems Perspective C.1: Node-level numbers for fleet reasoning**
>
> Fleet reasoning depends on a few node-level numbers that directly affect fleet design. The training state rule of about 14 bytes per parameter describes common Adam checkpoint state, with larger footprints when gradients, activations, metadata, or extra optimizer buffers are included. NVLink vs. high-bandwidth memory (HBM) bandwidth informs intra-node parallelism placement. Peak FLOP/s and HBM capacity feed model FLOPs utilization (MFU), effective FLOP/s, batch sizing, and model sharding. Table C.2 and the communication numbers below give inter-node and current-generation values. The memory hierarchy and roofline ridge points within a node (section A.4) provide the necessary baseline. Those node-level baselines make the fleet-level ratios easier to interpret.

> **Systems Perspective C.2: Three fleet numbers that matter most**
>
> - **Bandwidth gap**: NVLink per-direction bandwidth within a node is ~9× faster than an InfiniBand NDR port between nodes. This ratio determines where parallelism boundaries belong, with model parallelism within a node and data parallelism across nodes.
> - **MTBF scaling**: A cluster's mean time between failures is the single-component mean time to failure (MTTF) divided by the number of components ($1/N$ scaling). At 100,000 GPUs, expect a failure every 30 minutes.
> - **Effective utilization**: After MFU, scaling efficiency, and overhead losses compound, a cluster with 1,024 GPUs delivers roughly 34.0 percent of its peak FLOP/s as useful training work.

Back-of-envelope checks need only the condensed values in table C.2. Design and debugging work should use the detailed tables in each subsection.

Table C.2: Numbers Every Fleet Engineer Should Know (Quick Reference): One-page summary of the fleet-scale reference numbers in this section. See table C.3 and table C.4 for communication, table C.5 for compute, and table C.6 for coordination.

Category	Number	Use
Communication	NVLink per-direction bandwidth ~9× IB NDR	Parallelism boundary (in-node vs. cross-node)
Communication	IB NDR ~50 GB/s, ~5 µs one-way	Inter-node bandwidth and latency
Compute	MFU 30–50%, η_{scaling} ~83% @ 1K, ~56% @ 8K	Effective FLOP/s and scaling sanity checks
Coordination	MTBF 8K: ~366.2 min; 100K: ~30 min	Failure expectation and checkpoint cadence
Coordination	175B checkpoint ~2,450 GB (14 B/param common Adam state)	Recovery and storage sizing
Coordination	~77% of wall time left after overheads	Wall-clock utilization
Power & sustainability	AI rack ~70 kW; air limit ~30 kW	Cooling feasibility for the installed configuration
Power & sustainability	PUE liquid ~1.06, typical ~1.55; 10,000 H100s at 700 W ≈ 7 MW GPU-only proxy, before PUE	GPU-only proxy; add other IT loads before facility sizing (see section C.4)

Invariant scaling ratios

Three structural relationships govern how bandwidth, reliability, and collective communication behave as a cluster grows: the intra-to-inter-node bandwidth gap, the $1/N$ scaling of system MTBF, and the asymptotic per-GPU volume of Ring AllReduce. Because hardware ratios shift across generations, their numerical values must be re-evaluated for each target system.

Network hierarchy ratio

The bandwidth gap between intra-node and inter-node communication is an important architectural comparison. Chip-to-chip links (NVLink, ICI) connect through short, wide, dedicated paths on a shared substrate. Inter-node links (InfiniBand, Ethernet) must traverse cables, switches, and protocol stacks. These design differences often favor intra-node bandwidth, but lane counts, multiple network ports, and topology determine the actual ratio.

Currently, NVLink 4.0 provides 9× more per-direction bandwidth than an InfiniBand NDR port. The ratio compares one accelerator's aggregate intra-node connectivity with one network port; it is not a universal constant. Parallelism placement must compare the workload's traffic with the sustained bandwidth actually provisioned at each boundary.

System MTBF scaling

For N independent components under a constant-failure-rate assumption, system MTBF (equation C.1) scales inversely with component count:

$$\text{MTBF}_{\text{system}} = \frac{\text{MTTF}_{\text{component}}}{N} \tag{C.1}$$

Under this model the relation is exact rather than approximate. Doubling the cluster size halves the time between failures. At 100,000 GPUs with a per-GPU MTTF of 50,000 hours, the cluster experiences a GPU failure every 30 minutes. No fault tolerance strategy can avoid this; the question is how quickly the system recovers.

AllReduce overhead scaling

The bandwidth-optimal Ring AllReduce algorithm transfers $2(N_{\text{GPU}}-1)M/N_{\text{GPU}}$ bytes per participant, where M is the collective payload size and N_{GPU} is the number of participating GPUs (Patarasuk and Yuan 2009). As N_{GPU} grows large, this approaches $2M$ per GPU, so the per-GPU bandwidth term is nearly independent of the number of GPUs; aggregate cluster traffic still grows with the number of participants. This is why Ring AllReduce scales well in the bandwidth term. The latency term, however, grows as $2(N_{\text{GPU}}-1)\times\alpha$, making latency the bottleneck for small messages on large rings. This trade-off motivates hierarchical AllReduce strategies that use Ring AllReduce within nodes and Tree AllReduce across nodes.

Communication numbers

Whether a parallelism strategy can exchange activations on every layer or must amortize gradient updates across microbatches depends on which tier of the network hierarchy its traffic crosses. Each step down the hierarchy trades lower per-link bandwidth and higher one-way latency for larger physical reach across racks and data centers.

The bandwidth cliff between intra-node fabrics and scale-out network ports decides where tensor, pipeline, and data-parallel traffic must stay local (table C.3).

Table C.3: Communication Bandwidth Hierarchy: Per-direction bandwidth drops by roughly 9× crossing from intra-node (NVLink) to inter-node (InfiniBand). This ratio compares the H100 NVLink one-way rate against one InfiniBand NDR port and determines parallelism placement.

Interconnect	Bandwidth	Typical role
NVLink 4.0 (H100)	450 GB/s	Tensor/pipeline parallelism within a node
TPU v5p ICI	1,200 GB/s	Intra-pod model parallelism (Google)
PCIe Gen5 x16	64 GB/s	CPU-GPU data transfer, NIC attachment
IB GXDR (1.6 Tb/s)	200 GB/s	Next-gen inter-node (2026+)
IB XDR (800 Gb/s)	100 GB/s	Inter-node standard (2025+)
IB NDR (400 Gb/s)	50 GB/s	Current inter-node standard for AI clusters
IB HDR (200 Gb/s)	25 GB/s	Previous-gen inter-node
RoCE v2 (100 Gb/s)	12.5 GB/s	Budget clusters, inference fleets

One-way latency across interconnect tiers (table C.4) determines whether synchronous collective operations are viable. Whereas microsecond-scale fabrics support per-step gradient synchronization, wide-area links require asynchronous or amortized updates.

Table C.4: Communication Latency Hierarchy: Latency determines whether synchronous training is feasible. TCP/IP is roughly 10× slower than InfiniBand NDR, raising synchronization cost in latency-sensitive configurations. The cross-data-center row is the speed-of-light floor for an illustrative 8,000 km fiber path.

Interconnect	One-way latency	Implication
InfiniBand NDR	~5 µs	Low enough for synchronous AllReduce
InfiniBand HDR	~7 µs	Adequate for most training topologies
RoCE v2	~10 µs	Acceptable for data parallelism
TCP/IP (Ethernet)	~50 µs	Higher overhead; workload-dependent viability
Cross-data-center	~40,000 µs (40 ms)	Distance-dependent floor; amortize synchronization

These latency differences favor low-latency fabrics for frequent synchronization. Cross-data-center training must amortize transfers, place pipeline boundaries carefully, or relax synchronization when its learning algorithm permits.

Compute numbers

Raw peak FLOP/s is a necessary but misleading metric for fleet capacity planning. Single-device execution losses and additional scaling losses reduce effective throughput. Multiplying the two is valid only when their measurement windows exclude overlapping costs. Understanding these losses transforms fleet sizing from guesswork into engineering.

MFU measures what fraction of peak FLOP/s a training workload actually achieves. Well-optimized large-model training on current hardware achieves 30–50 percent MFU. The gap comes from memory stalls, kernel launch overhead, pipeline bubbles, and suboptimal operator fusion. These are workload-dependent reference values. Fleet-step MFU already includes in-step communication and pipeline losses; a separate scaling factor applies only to a single-device baseline that excludes those losses.

Scaling efficiency (η_{scaling}) measures how much useful computation survives as accelerators are added. Modeled scaling efficiency for the 175B reference job (table C.5) shows how per-GPU throughput erodes as cluster size grows at a fixed global batch.

Table C.5: Scaling Efficiency by Cluster Size: Modeled for the 175B reference job (eight-way tensor and eight-way pipeline parallelism, HDR InfiniBand, fixed GPT-3 batch); the 32-GPU row is a book convention. These values are not a universal function of cluster size, since workload, parallelism, batch, fabric, and baseline measurement determine the observed loss.

Cluster size	Scaling efficiency (η_{scaling})	Implication
32 GPUs	~90%	Near-linear regime; book convention, not modeled
256 GPUs	~87%	Tensor-parallel communication is the main loss
1,024 GPUs	~83%	Pipeline fill and drain become visible
8,192 GPUs	~56%	Pipeline bubble dominates at a fixed batch

Past a few thousand GPUs the fixed per-step costs of pipeline fill and drain dominate, which is why production systems grow the global batch, interleave pipeline stages, and overlap collectives with compute. Systems built this way, such as MegaScale, report near-linear scaling on thousands of GPUs (Z. Jiang et al. 2024).

Coordination numbers

At fleet scale, coordination (failure recovery, checkpointing, and maintenance) consumes a measurable fraction of wall-clock time. These numbers quantify the costs of keeping a large cluster running.

Failure rates by cluster size

Cluster MTBF scaling (table C.6, calculated for a per-GPU MTTF of 50,000 hours) dictates the required checkpoint frequency. As accelerator counts grow from hundreds to tens of thousands, the interval between expected hardware interruptions shrinks from days to minutes.

Table C.6: MTBF and Failure Probability by Cluster Size: GPU-only failure model with per-GPU MTTF of 50,000 hours. Real clusters also experience NIC, PSU, cable, and switch failures, so these GPU-only intervals overstate the true cluster MTBF. The probability column uses $\Pr(\geq 1 \text{ failure}) = 1 - e^{-t/\text{MTBF}}$ for $t = 24$ hours.

Cluster size	MTBF (GPU-only)	Minutes	Pr(**failure**) in 24 hours
256 GPUs	195.3 h	11,718.8 min	11.6%
2,048 GPUs	24.4 h	1,464.8 min	62.6%
8,192 GPUs	6.1 h	366.2 min	98%

Cluster size	MTBF (GPU-only)	Minutes	Pr(**failure**) in 24 hours
100,000 GPUs	0.50 h	30 min	100%

Failure is no longer an edge case at 8,192 GPUs and above. GPU-only MTBF is about 366.2 min at 8,192 GPUs, and the probability of at least one GPU failure within 24 hours is 98 percent. Fault tolerance is not optional at fleet scale; it is a prerequisite for completing any long-running training job. Chapter 7 covers the mechanisms in detail.

Checkpoint sizes

Checkpointing is the primary recovery mechanism, and its storage footprint grows linearly with model parameter count. A common mixed-precision Adam checkpoint serializes 14 bytes per parameter (2 bytes for BF16 weights and 12 bytes for FP32 master weights, momentum, and variance), which sets both total checkpoint capacity and write duration at a given storage bandwidth (table C.7). Gradients are normally transient and recomputed after restore rather than serialized as durable checkpoint state.

Table C.7: Checkpoint Sizes by Model Scale: Uses 14 bytes/parameter for common mixed-precision Adam checkpoint state. Write time assumes 100 GB/s aggregate storage bandwidth. Asynchronous checkpointing (chapter 7) can overlap writes with training, reducing the visible overhead.

Model size	Checkpoint size	Write time @ 100 GB/s
7B parameters	98 GB	0.98 s
70B parameters	980 GB	9.8 s
175B parameters	2,450 GB	24.5 s
1T parameters	14 TB	2.3 min

Overhead budgets

Failure recovery and checkpointing are only part of the wall-clock budget. At fleet scale, table C.8 separates four recurring categories of overhead that consume time not spent on useful training.

Table C.8: Overhead Budgets for Fleet-Scale Training: These are fractions of wall-clock time, and failure recovery is the largest of the four. Combined additively, the share of wall time left is $\approx 1.0 - (0.05 + 0.03 + 0.10 + 0.05) = 0.77 \approx 77\%$. Pipeline bubbles lie inside retained steps and are a program goodput loss, while the other three fall outside the steps, so this share is not a single goodput factor.

Overhead category	Typical budget	Lever
Pipeline bubbles	~5%	Increase microbatches per pipeline stage
Checkpointing	~3%	Async checkpointing, faster storage
Failure recovery	~10%	Faster detection, elastic rescheduling
Maintenance windows	~5%	Rolling upgrades, live migration

Power and sustainability numbers

Fleet-scale capacity planning and sustainability reporting require a few power numbers that every fleet engineer should know. At fleet scale, the key numbers are **rack power density**, **power usage effectiveness (PUE)**, and the **air-cooling limit**, which determine where construction is feasible and what the facility load will be (table C.9).

Table C.9: Power and Cooling Reference Numbers: Rack power densities, air-cooling limits, and PUE values that govern fleet-scale capacity planning and sustainability reporting. Each row pairs a typical value with the engineering decision it informs.

Quantity	Typical value	Use
Traditional rack	12 kW	Baseline for non-AI data centers
AI rack (current gen)	70 kW	Evaluate cooling capacity
AI rack (high-density)	100 kW	Direct-to-chip liquid
Air cooling limit	~30 kW per rack	Illustrative configuration limit; verify installed capacity
PUE (liquid-cooled AI)	~1.06	Best case: facility load ≈ IT load
PUE (best air-cooled)	~1.12	Hyperscale best practice
PUE (industry average)	~1.55	Sanity check for cost/carbon
H100 TDP	700 W per GPU	GPU-only proxy: 10,000 × 700 W = 7 MW

A GPU-only power proxy (MW) = (number of GPUs × thermal design power (TDP) per GPU)/10^6 excludes host CPUs, memory, storage, and networking. Facility load = total IT load × PUE; applying PUE to the GPU-only proxy omits those other IT loads and is not a complete facility budget. A 10,000-GPU H100 cluster at 700 W each gives a 7 MW GPU-only proxy; multiplying by PUE 1.55 gives 10.8 MW before accounting for the other IT loads. Section C.4 carries this power arithmetic into facility overhead, and chapter 15 converts it into carbon and cost.

Current hardware reference (c. 2025–2026)

Frontier accelerators from the 2025–2026 generation provide the point-in-time per-chip specifications needed for back-of-envelope fleet sizing (table C.10), illustrating how compute throughput, memory capacity, and interconnect bandwidth evolve at different rates across vendors.

Table C.10: Fleet-Scale Hardware Reference (c. 2025–2026): Per-accelerator specifications for the 2025–2026 generations. The B200, MI300X, and TPU v6e (Trillium) represent the frontier of fleet-scale compute density and memory bandwidth. Intra-node link rates are bidirectional aggregates per accelerator. Google does not publish a per-chip TDP for the TPU v6e.

Spec	NVIDIA B200	AMD MI300X	Google TPU v6e
BF16/FP16 Peak	2,250 TFLOP/s	1,307 TFLOP/s	918 TFLOP/s
Memory Bandwidth	8 TB/s	5.30 TB/s	1.64 TB/s
HBM Capacity	180 GB	192 GB	32 GB
Intra-Node Link	1,800 GB/s (NVLink 5.0)	896 GB/s (Infinity Fabric) (AMD 2023)	800 GB/s (ICI) (Google Cloud 2026b)
TDP	1,000 W	750 W	Not published

At the cluster level, a single DGX H100 SuperPOD scalable unit packages 32 DGX H100 nodes (256 H100 GPUs) (NVIDIA 2023), whereas full data-center deployments link dozens of such pods together. Meta's Grand Teton clusters reach 24,576 H100 GPUs (Meta Engineering 2024), Google's Tensor Processing Unit (TPU) v5p pods scale to 8,960 chips, and the largest announced clusters as of 2025 exceed 100,000 accelerators.

C.3 Scaling Physics

When a training job expands from a single 8-GPU server to thousands of accelerators across a fabric, network transfers, synchronization barriers, and pipeline bubbles consume a growing share of each step. The reference numbers in section C.2 describe what individual links and devices can deliver; scaling physics explains how much of that capacity survives as devices are added. Reasoning starts from the ceiling on a single accelerator (the roofline) and extends it to the cluster through three models: Amdahl's Law for fleet serial overhead, the communication-computation ratio, and weak scaling.

C.3.1 From the accelerator roofline to the fleet

Fleet performance rests on per-accelerator performance. Section A.3.1 states the single-accelerator roofline, $R_{\text{attain}} = \min(R_{\text{peak}}, I \times \text{BW})$, and its ridge point $I_{\text{ridge}} = R_{\text{peak}}/\text{BW}$, which is about 295.2 FLOP/byte for an H100. A kernel below the ridge is memory bound, and a kernel above it is compute bound.

The roofline is a single-accelerator bound, but its logic is what scales. At fleet level the binding bandwidth is no longer HBM but the interconnect, and the bytes moved are the gradients and activations crossing the network. The ratio of network time to local arithmetic time is the communication-computation ratio (section C.3.3), the roofline of the fleet and the lens the rest of this section applies.

On a real cluster, that shift from HBM bandwidth to network and synchronization limits shows up directly in step throughput: when a training job provisioned with 4,096 GPUs achieves only 2.5× the throughput of a 1,024-GPU run, the implied scaling efficiency is about 51.9 percent at the larger size. The three subsections that follow separate whether that loss comes from an unavoidable Amdahl serial fraction, an unfavorable communication-computation ratio, or a fixed-problem strong-scaling regime.

C.3.2 Amdahl's Law at fleet scale

When a distributed training loop spends 10 percent of every step waiting on exposed AllReduce collectives, pipeline bubbles, or checkpoint flushes, adding thousands of extra GPUs cannot accelerate the fixed job by more than $10\times$. Amdahl's Law establishes that the maximum speedup of a parallel system is limited by its serial fraction r. At fleet scale, the "serial fraction" includes more than sequential code; it covers every operation that forces all N GPUs to wait:

- **AllReduce synchronization**: All GPUs must complete their gradient computation before any can proceed.
- **Pipeline bubbles**: The warmup and cooldown phases of pipeline parallelism leave stages idle.
- **Checkpoint writes**: Even asynchronous checkpoints contend for storage bandwidth.
- **Python-level overhead**: Single-threaded operations in the training loop (data loading, metric logging).

To see the fleet-scale implications, consider a training workload where 10 percent of wall-clock time is spent in synchronization, communication, and other serial overhead. Amdahl's Law gives the following speedups:

- With 32 GPUs: 7.8× speedup (~24 percent scaling efficiency)
- With 256 GPUs: 9.7× speedup (diminishing returns begin)
- With 1,024 GPUs: 9.9× speedup (approaching the ceiling)
- With 8,192 GPUs: 10× speedup (nearly at the Amdahl limit)
- With $N \to \infty$: capped at 10×

With 10 percent serial overhead, no amount of hardware can deliver more than 10× speedup on this fixed workload. This is why fleet-scale training does not add more GPUs to the same problem. It scales the problem (weak scaling) to keep the serial fraction small relative to the total work.

Systems Perspective C.3: The compound overhead trap

The 10 percent serial fraction above is optimistic. In practice, fleet-scale serial overhead accumulates from multiple sources: 5 percent pipeline bubbles + 3 percent checkpointing + 10 percent failure recovery + 5 percent maintenance. These are not all strictly serial in the Amdahl sense (some overlap with computation), but they illustrate how quickly small per-category overheads compound into significant throughput loss.

C.3.3 The communication-computation ratio

When 256 GPUs train a 350M-parameter model over InfiniBand, gradient AllReduce takes three times longer than the forward and backward passes combined, leaving the accelerators idle for most of the step. Any distributed training strategy must determine whether computation between synchronization points takes long enough to hide the communication. Equation C.2 defines the **communication-computation ratio** (ρ) that answers this directly:

$$\rho = \frac{T_{\text{comm}}(N)}{T_{\text{compute}}/N} \tag{C.2}$$

When $\rho < 1$, computation takes longer than communication; hiding that communication still requires independent work and sufficient resources. When $\rho > 1$, the system is communication bound, and GPUs spend more time waiting for data than computing on it.

The ratio is the time-domain form of the **communication intensity** (CI) defined in section 1.2.4, the bytes a step moves across the network per FLOP it executes locally. If each device moves M bytes across the network per step at effective network bandwidth $\text{BW}_{\text{network}}$ and executes its O local FLOPs at realized rate R, then in the bandwidth-dominated regime

$$\rho \approx \frac{M/\text{BW}_{\text{network}}}{O/R} = \text{CI} \times \frac{R}{\text{BW}_{\text{network}}}, \qquad \text{CI} = \frac{M}{O} \tag{C.3}$$

Equation C.3 shows that exposed communication begins when CI exceeds $\text{BW}_{\text{network}}/R$, the network counterpart of the roofline ridge point. Because a device executes far more FLOP/s than its network link moves bytes per second, that threshold is a small fraction of a byte per FLOP, and a faster accelerator lowers it further. An accelerator upgrade alone can therefore turn a compute-bound job into a communication-bound one.

Communication-to-computation ratios across representative scenarios (table C.11) reveal why the parallelism strategy must match the workload's arithmetic footprint. Comparing a multi-billion-parameter model against a smaller architecture on the same interconnect shows how quickly collective synchronization can overtake local compute.

Table C.11: Communication-Computation Ratio (ρ): When $\rho \ll 1$, communication can potentially be hidden if dependencies permit. When $\rho \gg 1$, the system is communication-bound and GPUs sit idle waiting for data. The NVLink row is a bandwidth-only point-to-point lower bound; a tensor-parallel collective also incurs topology, latency, and contention costs.

Scenario	$T_{\text{comm}}(N)$	T_{compute}/N	ρ
Data parallel, 7B model, 256 GPUs	560.4 ms (AllReduce over IB NDR)	~5 s (fwd+bwd)	0.11
Data parallel, 340M model, 256 GPUs	29.6 ms (AllReduce over IB NDR)	~10 ms (fwd+bwd)	3
Ideal one-way 16 MB transfer (NVLink)	~0.04 ms (bandwidth-only lower bound)	~1 ms (one layer)	0.036

The 7B model on 256 GPUs achieves $\rho = 0.11$, so communication takes about 11.2 percent as long as computation, which can be partially overlapped. The 340M model on the same cluster has $\rho = 3$. Communication dominates, which makes this configuration communication-bound. The solution is either to use fewer GPUs (reduce N in the AllReduce) or to increase the computation per step

(larger batch size, gradient accumulation). For the single 16 MB point-to-point payload, the ideal bandwidth-only lower bound is $\rho = 0.036$, which is not a tensor-parallel collective model and cannot by itself establish whether an implementation is compute-bound.

C.3.4 Weak scaling at fleet scale

Provisioning 8,192 GPUs to finish a fixed 7B-parameter training run faster quickly hits Amdahl's synchronization wall. Using those same 8,192 GPUs to train a 70B- or 700B-parameter model keeps per-device arithmetic work high enough to amortize communication. Amdahl's Law paints a pessimistic picture because it assumes a fixed problem size. In practice, fleet-scale training often follows **weak scaling**: the problem size (tokens, data, model parameters) grows proportionally with the number of GPUs. Gustafson's Law captures this more optimistic view (Gustafson 1988).

Weak scaling reflects how fleets are actually used, not merely a mathematical convenience. Engineers do not use 8,192 GPUs to train a 7B model faster; they use them to train a 70B or 700B model in reasonable time. As models and training batches grow, engineers often increase the compute performed between synchronization points. For dense Transformers, compute is commonly estimated as about 6× parameters × tokens, while gradient communication scales with parameter bytes. The communication-computation ratio ρ therefore depends on batch size, sequence length, accumulation, and parallelism layout rather than on parameter count alone.

Systems Perspective C.4: The compound loss of fleet utilization

A 1,024-GPU H100 cluster has a peak aggregate throughput of 1,012,736 TFLOP/s. This illustrative decomposition keeps the factors disjoint: MFU represents per-device kernel and memory efficiency, scaling efficiency represents distributed communication and pipeline losses, and runtime goodput represents checkpointing, failure recovery, and maintenance. The effective throughput is:

$$R_{\text{eff}} = N R_{\text{peak}} \times \text{MFU} \times \eta_{\text{scaling}} \times \eta_{\text{runtime}}$$

$$= 1,012,736 TFLOP/s \times 0.50 \times 0.83 \times 0.82 \approx 344,634.1 \text{ TFLOP/s}$$

The cluster delivers 34.0 percent of its peak FLOP/s as useful training work. The remaining 66.0 percent is explained by retained factors of 50 percent MFU, 83 percent scaling efficiency, and 82 percent runtime goodput.

These losses reflect the physics of fleet-scale computation rather than a failure of engineering. Additional GPUs can increase useful throughput while reducing efficiency, but the gain must be measured against smaller configurations. The goal is not to reach 100 percent utilization; the goal is to deliver trained models faster than any smaller configuration could.

C.4 Thermal and Power Physics

A cluster design calling for 10,000 H100 GPUs at 700 W each represents a 7 MW GPU-only power proxy and draws 10.8 MW of facility power at a typical PUE of 1.55, the electrical load of a small town even before adding host CPUs, memory, storage, and networking inside the IT boundary. Because every watt of electrical power delivered to silicon converts to heat that must be removed, thermal and power limits govern where a cluster can be built, how densely accelerators can be packed into each rack, and what the facility will cost to operate. A cluster that is architecturally sound in software and networking but exceeds its power-delivery or heat-removal envelope cannot be deployed.

C.4.1 Power density wall

The shift from traditional data center workloads to AI training has multiplied the power each rack draws, and table C.12 quantifies the gap. As accelerator TDPs climb toward a kilowatt per chip, standard air-cooled server rooms quickly run out of thermal headroom.

Table C.12: Power Density: Traditional vs. AI Workloads: AI racks consume roughly 5.8× the power of traditional racks. The ~30 kW air-cooling value is a configuration reference, not a universal physical ceiling. Higher rack density requires checking the installed airflow, heat exchangers, and server cooling design.

Configuration	Power per rack	Cooling implication
Traditional data center	12 kW	Standard air cooling sufficient
AI cluster (current gen)	70 kW	Evaluate cooling capacity
AI cluster (high-density)	100 kW	Evaluate direct-to-chip cooling
Air cooling limit	~30 kW	Illustrative forced-air configuration limit

The 5.8× increase in rack power density between traditional and AI workloads is a binding constraint on facility design, not merely an engineering challenge. Existing data centers built for 12 kW racks cannot be "retrofitted" with AI hardware. The electrical distribution, cooling infrastructure, and floor loading must all be redesigned. This is why new AI-focused data centers are being built from the ground up with liquid cooling as the baseline assumption.

C.4.2 The energy hierarchy at scale

In a data center at the 2023 industry-average PUE, delivering 1 MW of electrical power to GPU racks requires another 580 kW to run the chillers, fans, and power conversion equipment that keep the room operating. PUE measures facility energy overhead: total facility energy divided by IT equipment energy (The Green Grid 2007; Uptime Institute 2022). A PUE of 1.0 would mean all facility energy is delivered to IT equipment; values above 1.0 represent overhead for cooling, power conversion, lighting, and other facility systems. Cooling technology drives much of that overhead, so PUE falls from legacy enterprise facilities to liquid-cooled AI data centers (table C.13).

Table C.13: Power usage effectiveness by facility type: Lower is better. The listed PUE values are facility scenarios. Direct liquid cooling can reduce heat-removal overhead, but PUE also depends on climate, heat rejection, power conversion, and load. The gap between the 2023 industry average (1.58) and liquid-cooled (1.06) represents about a 32.9 percent reduction in total facility power at fixed IT load; the non-IT overhead drops from 0.58 MW to 0.06 MW per MW of IT load.

Data center type	PUE	Overhead per MW of IT load
Liquid-cooled AI data center	1.06	60 kW cooling + infrastructure
Best-in-class air-cooled	1.12	120 kW overhead
Industry average (2022)	1.55	550 kW overhead
Industry average (2023)	1.58	580 kW overhead

For fleet-scale cost calculations, PUE directly multiplies the electricity bill. A 10 MW IT load in a data center at the 2023 industry average (PUE 1.58) requires 15.8 MW total. The same load in a liquid-cooled facility (PUE 1.06) requires only 10.6 MW, a savings of 5.2 MW. For a facility that runs continuously at the reference on-premises electricity price of $0.07/kWh, that difference is worth about $3.2M per year. Chapter 15 covers the full sustainability implications.

Summary

Fleet scale changes every constraint a single machine imposes. ML fleets combine the tight coupling of high-performance computing for training with the elastic fault tolerance of warehouse-scale computing for serving, and they cannot adopt either tradition wholesale. NVLink provides about 9× more per-direction bandwidth than an InfiniBand NDR port, which is why bandwidth-hungry parallelism stays inside the node. Cluster MTBF scales as $1/N$, so at 8,192 GPUs a GPU fails about every 366.2 minutes, with 98 percent probability of at least one failure in 24 hours. After disjoint factors for MFU (about 50 percent), scaling efficiency (about 83 percent), and runtime goodput (about 82 percent), a 1,024-GPU cluster delivers about 34.0 percent of peak FLOP/s as useful work.

Three considerations govern the design response. The communication-computation ratio ρ indicates the potential for overlap and the pressure from communication; dependencies and available resources determine how much communication can be hidden behind computation. Weak scaling, which grows the problem with the fleet, preserves useful work per device, while strong scaling remains the right choice when the goal is to finish fixed work sooner. Rack power density, about 5.8× the traditional reference for current AI racks, determines the cooling design, subject to the installed heat-removal capacity rather than a universal threshold.

D

Communication Foundations

Distributed training spends a substantial fraction of execution time synchronizing gradients, redistributing activations, and routing tokens across network fabrics. Rather than cluttering chapter text with step-by-step proofs, this appendix provides the quantitative communication cost models used throughout the book: point-to-point formulations (α-β and LogGP), collective execution bounds (AllReduce, AllGather, ReduceScatter, AllToAll), pipeline bubble math, and gradient compression thresholds.

How to Use This Appendix

This appendix is designed as a reference when translating a distributed training symptom ("AllReduce is slow," "pipeline bubbles are killing throughput," "should we compress gradients?") into a quantitative diagnosis. Conventions used here follow the book-wide notation. α denotes per-message startup latency, and β denotes link bandwidth in bytes per second. n denotes point-to-point message size, and n^* denotes the α-β crossover point. N denotes the number of participating GPUs, and M denotes collective payload size in bytes.

- **When AllReduce is slow**: Use section D.2 to compute the expected time and compare ring vs. tree costs.
- **When choosing ring vs. tree**: Use section D.3 and the crossover analysis to match algorithm to message size.
- **When pipeline bubbles dominate**: Use section D.4 to compute bubble fraction and determine the required microbatch count.
- **When gradient compression is proposed**: Use section D.5 to run the break-even calculation before committing engineering effort.

D.1 The α-β Communication Model

> **Systems Perspective D.1: Why this matters**
>
> The time a gradient synchronization takes across 256 GPUs must be estimated before a cluster topology is chosen. The α-β model plays the role of a roofline for networks, giving a first-order prediction of transfer time from two hardware parameters and the message size.

Every point-to-point transfer on a network follows the same pattern: the sender pays a fixed startup cost to initiate the message, then transmits the payload at a rate determined by the link's bandwidth. This two-parameter model captures the essential physics of network communication. Specifically, the **α-β model** (equation D.1) predicts the time to transfer a point-to-point message of n bytes:

$$T(n) = \alpha + \frac{n}{\beta} \tag{D.1}$$

Here α is the **startup latency** (seconds), the fixed cost of initiating a message including software overhead, network interface card (NIC) processing, and routing setup. β is the **link bandwidth** (bytes/s), the sustained data rate after startup. n is the **message size** (bytes). The model decomposes every transfer into exactly two costs: a per-message cost (α) paid regardless of size, and a per-byte

cost (n/β) that scales linearly with the payload. This decomposition drives all the algorithm selection decisions in section D.3, and the baseline α and β parameters span more than an order of magnitude across intra-node NVLink and PCIe, inter-node InfiniBand and RoCE fabrics, and kernel TCP stacks (table D.1).

Table D.1: α-β Parameters by Interconnect: Effective startup latency and sustained bandwidth for interconnects used in ML clusters. α follows the convention of table 5.4 in chapter 5, the effective per-message launch latency including the communication library's software stack, which is why it reaches microseconds even on NVLink. The InfiniBand NDR entry is the single anchor that every worked example in this appendix uses, and it lies inside that table's NDR range. The β column uses one-way or per-direction bandwidth, so the H100 NVLink entry is half of the 900 GB/s bidirectional aggregate specification. NVLink and PCIe operate within a node; InfiniBand and RoCE operate between nodes. TCP latency is dominated by kernel software stack overhead.

Interconnect	α (latency)	β (bandwidth)	Typical role
NVLink 4.0 (H100)	1–2 µs	450 GB/s	Intra-node GPU-to-GPU
PCIe Gen5	2–5 µs	64 GB/s	CPU-GPU, NIC-GPU
InfiniBand NDR (400 Gb/s)	1.5 µs	50 GB/s	Inter-node (high-end clusters)
InfiniBand HDR (200 Gb/s)	2–5 µs	25 GB/s	Inter-node (previous generation)
RoCE v2 (100 Gb/s)	5–10 µs	12.5 GB/s	Inter-node (Ethernet clusters)
TCP/IP (Ethernet)	50 µs	Varies	Control plane, non-RDMA fallback

D.1.1 Latency-dominated vs. bandwidth-dominated regimes

Sending a 64-byte control flag over InfiniBand NDR takes nearly the same five microseconds as sending a 64 KB tensor, whereas a 100 MB gradient shard spends almost all of its transfer time streaming bytes at line rate. That contrast separates the α-β model into two operating regimes: for small messages, the startup latency α dominates and the link spends most of its time initiating transfers, whereas for large messages, the bandwidth term n/β dominates and the startup cost becomes negligible. The **crossover message size** n^*, defined by equation D.2, is the point where both terms contribute equally:

$$n^* = \alpha \times \beta \tag{D.2}$$

Below n^*, communication is **latency dominated**: sending more small messages wastes time on repeated startups. Above n^*, communication is **bandwidth dominated**: payload transfer dominates; reducing bytes or increasing effective bandwidth can lower its cost.

Napkin Math D.1: Worked example: Crossover on InfiniBand NDR

Setup: InfiniBand NDR has $\alpha = 1.5$ µs and $\beta = 50$ GB/s.

Question: At what message size does the bandwidth term equal the latency term?

Math:

$n^* = \alpha \times \beta = 1.5\ \mu\text{s} \times 50\ \text{GB/s} = 75\ \text{KB}$

Analysis: Messages smaller than 75 KB are latency dominated on IB NDR. Gradient tensors from a single layer of a large model are typically 10–100 MB, well above this threshold. Control messages, heartbeats, and barrier synchronizations, however, are well below it.

Systems insight: This is why NCCL and Gloo batch small gradient tensors into larger "buckets" before launching AllReduce. Sending each tensor individually would pay the 1.5 µs startup cost hundreds of times per iteration instead of a handful.

D.1.2 The LogGP model extension

The α-β model assumes a message is a single, atomic transfer. In practice, networks pipeline messages: a sender can inject a new packet before the previous one has been fully received. The **LogGP model** (Alexandrov et al. 1995) extends LogP with a long-message gap parameter:

- g **(gap)**: The minimum interval between consecutive message injections at the sender. This models the NIC's injection rate limit.
- o_{send}, o_{recv} **(overhead)**: The CPU time consumed by the processor to initiate or complete a message, during which it cannot do useful compute.

For a long message of n bytes, LogGP models transfer time by separating network transit from endpoint processor overhead and per-byte injection gaps:

$$T_{\text{long}}(n) \approx o_{\text{send}} + L_{\text{LogGP}} + (n-1)G + o_{\text{recv}}$$

Here L_{LogGP} is the network latency, o_{send} and o_{recv} are sender and receiver overheads, G is the gap per byte for long messages, and g controls the minimum interval between consecutive message injections. Informal alpha-beta-style approximations can be useful for ML profiling, but the standard LogGP model keeps these gap parameters distinct from a simple bandwidth term.

In most ML training scenarios, the α-β model is sufficient because gradient tensors are large enough that pipelining effects are secondary to raw bandwidth. The LogGP model becomes important when analyzing the overhead of many small control messages or when modeling the interaction between communication and computation overlap, situations that arise in fine-grained pipeline parallelism and asynchronous gradient updates. Whereas the α-β and LogGP equations describe single point-to-point transfers, collective operations in which all GPUs participate compose multiple point-to-point transfers, and their cost formulas derive directly from the model above.

D.2 Collective Operation Complexity

> Systems Perspective D.2: Why this matters
>
> Every distributed training step ends with a collective synchronization. The cost of that collective, rather than the cost of the matrix multiplies, often determines whether scaling from 8 GPUs to 256 is feasible. Understanding the exact complexity of each collective reveals which parallelism strategy (data, tensor, pipeline, expert) will dominate wall time at a given scale.

When N GPUs synchronize gradients, reconstruct sharded weights, or route tokens at the end of a computation phase, the collective primitive coordinates a multi-step schedule of point-to-point transfers across every participating rank. Each collective therefore has a characteristic bandwidth term (how many bytes traverse the bottleneck link) and latency term (how many sequential message steps must complete). The formulas below use the **bandwidth-optimal ring algorithm** unless otherwise noted; section D.3 discusses when other algorithms are preferable.

D.2.1 AllReduce

AllReduce is the workhorse of data-parallel training: every GPU starts with a local gradient tensor of size M bytes and ends with the globally reduced (summed) result. The ring algorithm decomposes AllReduce into a ReduceScatter phase followed by an AllGather phase. Ring AllReduce time is given by equation D.3:

$$T_{\text{ring}} = 2 \cdot \frac{N-1}{N} \cdot \frac{M}{\beta} + 2(N-1) \cdot \alpha \tag{D.3}$$

The factor $2(N-1)/N$ approaches 2 as N grows. Each rank sends that factor times M bytes and receives the same amount. With full-duplex links, sending and receiving overlap; β is the effective bandwidth in one direction, so sent-plus-received bytes must not be divided by that one-direction rate. The latency term grows linearly with N, which makes the ring algorithm expensive in latency for large GPU counts.

Tree AllReduce trades that linear latency growth for logarithmic depth. Equation D.4 charges both the bandwidth and the latency term once per tree level, $2\log_2 N$ times in total:

$$T_{\text{tree}} = 2\log_2 N \cdot \frac{M}{\beta} + 2\log_2 N \cdot \alpha \tag{D.4}$$

This unsegmented tree algorithm trades bandwidth efficiency for latency efficiency: the latency term grows as $\log_2 N$ instead of N, but the bandwidth term sends the full message M at each tree level instead of M/N chunks. Comparing the two formulas shows why large payloads favor ring topologies while small control synchronizations favor trees.

Napkin Math D.2: Worked example: Ring vs. tree AllReduce

Setup: 256 H100 GPUs connected via InfiniBand NDR (α = 1.5 µs, β = 50 GB/s). The gradient tensor is 1 GB (a 500M-parameter model in FP16).

Question: How long does AllReduce take with ring vs. tree?

Ring AllReduce:

$T_{\text{ring}} \approx 40.6$ ms

Tree AllReduce:

$T_{\text{tree}} \approx 320$ ms

Analysis: For this 1 GB message, ring AllReduce takes 40.6 ms, and tree takes 320 ms. The ring is significantly faster because it distributes the bandwidth load across all links. The tree's logarithmic latency advantage is negligible compared to its bandwidth penalty for large messages.

Systems insight: For the large gradient tensors typical of data-parallel training, ring (or ring-based) algorithms dominate. This unsegmented tree model favors small messages. Production tree variants can pipeline or partition messages and must be evaluated with their own cost models.

D.2.2 AllGather

AllGather is the complement of ReduceScatter: each GPU starts with a M_{coll}/N-sized shard and ends with the complete M_{coll}-byte tensor. It is the core communication primitive in fully sharded data parallelism (FSDP) for reconstructing parameters before each forward pass, and in tensor parallelism for reassembling split activations. Ring AllGather execution time (equation D.5) requires half the transfer volume of AllReduce:

$$T_{\text{allgather}} = \frac{N-1}{N} \cdot \frac{M_{\text{coll}}}{\beta} + (N-1) \cdot \alpha \tag{D.5}$$

For this comparison, the full gathered tensor has the same byte size as each rank's AllReduce input payload. In this model, AllGather takes half the communication time of AllReduce for that matched tensor size and link parameters: it uses one ring pass instead of two. Local reduction work and implementation overheads can change the measured ratio.

D.2.3 ReduceScatter

ReduceScatter is the dual of AllGather: each GPU starts with a full M_{coll}-byte tensor and ends with a M_{coll}/N-sized shard that contains the globally reduced values for that shard. It is the gradient synchronization primitive in FSDP and ZeRO, replacing the full AllReduce with a cheaper operation when each GPU only needs its own parameter shard's gradient. Ring ReduceScatter execution time (equation D.6) provides the gradient sharding primitive:

$$T_{\text{reducescatter}} = \frac{N-1}{N} \cdot \frac{M_{\text{coll}}}{\beta} + (N-1) \cdot \alpha \tag{D.6}$$

The symmetry with equation D.5 is not a coincidence: ring AllReduce decomposes into one ReduceScatter followed by one AllGather, and the costs add up exactly to equation D.3. Because

each phase traverses the ring once, sharded data-parallel schemes that defer the AllGather until the next forward pass cut immediate post-backward communication in half.

D.2.4 AllToAll

AllToAll is the most general collective: each GPU sends a distinct M/N-sized chunk to every other GPU, and receives a distinct chunk from each. It is the routing primitive for mixture-of-experts (MoE) architectures, where tokens must be dispatched to the correct expert and the results gathered back. Pairwise-exchange AllToAll execution time (equation D.7) trades personalized data chunks across peers:

$$T_{\text{alltoall}} = \frac{N-1}{N} \cdot \frac{M}{\beta} + (N-1) \cdot \alpha \tag{D.7}$$

Although the formula matches AllGather and ReduceScatter in form, the communication pattern is structurally different. AllToAll is a personalized exchange in which each GPU sends different data to each peer, which makes it harder to overlap with computation and much more sensitive to network bisection congestion.

D.2.5 Broadcast

Broadcast sends a single M-byte tensor from one root GPU to all $N-1$ others. It is used for distributing initial model weights, updated learning rates, and configuration changes during training. Tree Broadcast execution time (equation D.8) scales logarithmically with the participating GPU count:

$$T_{\text{broadcast}} = \log_2 N \left(\alpha + \frac{M}{\beta} \right) \tag{D.8}$$

Broadcast is structurally asymmetric (one sender, many receivers), which makes a tree topology natural. In an unsegmented tree broadcast, the full M-byte payload is forwarded at each tree level, so both the latency and bandwidth terms appear along the $\log_2 N$-level critical path. Segmented or scatter-allgather broadcast variants can reduce the bandwidth term, but they require a different algorithm-specific model.

Collective complexity summary

With the individual formulas in place, table D.2 provides a compact reference for comparing the cost structure of each collective operation. Checking the bandwidth and latency columns side by side shows how each collective scales as either payload size or device count grows.

Table D.2: Collective Operation Cost Summary: Bandwidth and latency terms for ring, tree, and pairwise-exchange algorithms. The bandwidth term determines cost for large messages (gradients, parameters); the latency term determines cost for small messages (scalars, barriers). N = GPU count, M = per-rank input payload in bytes, and M_{coll} = full tensor size in bytes for AllGather and ReduceScatter.

Operation	Bandwidth term	Latency term	Primary use case
AllReduce (Ring)	$2 \cdot \frac{N-1}{N} \cdot \frac{M}{\beta}$	$2(N-1) \cdot \alpha$	Data-parallel gradient sync
AllReduce (Tree)	$2\log_2 N \cdot \frac{M}{\beta}$	$2\log_2 N \cdot \alpha$	Small-message sync
AllGather (Ring)	$\frac{N-1}{N} \cdot \frac{M_{\text{coll}}}{\beta}$	$(N-1) \cdot \alpha$	FSDP parameter reconstruction
ReduceScatter (Ring)	$\frac{N-1}{N} \cdot \frac{M_{\text{coll}}}{\beta}$	$(N-1) \cdot \alpha$	FSDP/ZeRO gradient sharding
AllToAll (Pairwise)	$\frac{N-1}{N} \cdot \frac{M}{\beta}$	$(N-1) \cdot \alpha$	MoE expert routing
Broadcast (Tree)	$\log_2 N \cdot \frac{M}{\beta}$	$\log_2 N \cdot \alpha$	Weight distribution, config sync

Two patterns emerge from this table. First, the ring and pairwise-exchange models contain the $(N-1)/N$ factor, which approaches 1 for large N. Per-rank communication therefore approaches a constant multiple of the relevant size: about $2M$ for ring AllReduce, about M_{coll} for ring AllGather or ReduceScatter, and about M for pairwise AllToAll. Aggregate traffic still scales with the number of ranks. Second, the latency term grows linearly with N for ring algorithms but logarithmically for tree algorithms, creating the algorithm selection trade-off analyzed next.

D.3 Algorithm Selection

> Systems Perspective D.3: Why this matters
>
> NCCL automatically selects a collective algorithm based on message size and GPU topology, but understanding why it chooses ring for large messages and tree for small ones helps diagnose when the auto-selection is suboptimal and how to structure communication to stay in the efficient regime.

D.3.1 Ring vs. tree vs. recursive halving-doubling

Sending a 1 GB gradient tensor across 256 GPUs favors an algorithm that spreads bytes across every link, whereas synchronizing a 4-byte scalar loss flag favors an algorithm that finishes in the fewest hop steps. The ring algorithm minimizes bandwidth usage but pays a latency cost that grows linearly with N. The tree algorithm minimizes latency but wastes bandwidth because it sends the full message at each level. Recursive halving-doubling splits the difference: it uses halving (like a tree) for the reduce phase and doubling for the gather phase, achieving near-optimal bandwidth with $\mathcal{O}(\log N)$ latency steps.

The decision boundary between ring and tree depends on both message size and participant count. For the power-of-two case modeled by equation D.3 and equation D.4, equating the two times gives the crossover in equation D.9, which table D.3 translates into practical selection rules across payload regimes:

$$M_{\text{crossover}}(N) = n^* \times \frac{(N-1) - \log_2 N}{\log_2 N - (N-1)/N}, \qquad n^* = \alpha\beta \tag{D.9}$$

- $M > M_{\textbf{crossover}}(N)$ (large gradients): The ring has the lower modeled cost because its smaller bandwidth coefficient outweighs its additional startup steps.
- $M < M_{\textbf{crossover}}(N)$ (barriers, scalars): The tree has the lower modeled cost because its logarithmic startup term outweighs its additional bandwidth cost.
- $M \approx M_{\textbf{crossover}}(N)$ (medium tensors): The two modeled costs are similar, so benchmark them alongside recursive halving-doubling on the target topology.

Table D.3: Algorithm Selection Guide: Strengths and weaknesses of ring, tree, and recursive halving-doubling relative to the participant-dependent ring-tree crossover $M_{\text{crossover}}(N)$ in equation D.9.

Regime	Ring	Tree	Recursive halving-doubling
$M > M_{\textbf{crossover}}(N)$	Lower cost under the displayed models	Higher bandwidth term	Candidate with near-optimal bandwidth and logarithmic steps
$M < M_{\textbf{crossover}}(N)$	Higher linear startup term	Lower cost under the displayed models	Candidate with logarithmic startup steps
$M \approx M_{\textbf{crossover}}(N)$	Benchmark on the target topology	Benchmark on the target topology	Benchmark on the target topology
Power-of-2 N	Works for any N	Matches the displayed $\log_2 N$ model	Requires $N = 2^m$

D.3.2 Hierarchical AllReduce

Modern clusters have a two-level topology: GPUs within a node are connected by NVLink (450 GB/s), while nodes are connected by InfiniBand (50 GB/s). Hierarchical AllReduce exploits this asymmetry in three phases:

1. **Intra-node reduce** (NVLink): Each node reduces its 8 local GPUs to a single result using NVLink's 450 GB/s bandwidth. This is fast because NVLink's per-direction rate is 9× that of an InfiniBand NDR port.
2. **Inter-node AllReduce** (InfiniBand): An AllReduce combines the per-node results across nodes. Only one representative per 8 local GPUs participates, so the inter-node ring has one-eighth as many participants as a flat GPU-level ring.
3. **Intra-node broadcast** (NVLink): Each node broadcasts the global result to its 8 local GPUs.

The hierarchical approach reduces the inter-node ring size from all GPUs to one representative per 8 local GPUs, cutting the latency-step term by about 8× while keeping local aggregation on fast intra-node links. The inter-node representative still exchanges the full reduced tensor; the gain is not a universal reduction of payload bytes by the number of local GPUs. NCCL selects among topology-aware algorithms such as Ring, Tree, CollNet variants, NVLS/NVLSTree, and PAT based on the collective, message size, topology, architecture, and version. Hierarchical communication is common on multi-node systems but is not a universal default. The cost structure of collectives lets an engineer reason about data-parallel overhead. Pipeline parallelism, in contrast, introduces a different kind of overhead, idle time caused by the sequential dependency between pipeline stages.

D.4 Pipeline Arithmetic

> Systems Perspective D.4: Why this matters
>
> Pipeline parallelism splits a model across multiple GPUs so that each GPU holds only a fraction of the layers. The trade-off is **bubble time**: idle cycles where GPUs wait for activations or gradients from upstream or downstream stages. Quantifying the bubble fraction reveals whether pipeline parallelism is worth the complexity at a given scale.

D.4.1 Bubble fraction

When a model is split across eight pipeline stages, Stage 8 sits idle while Stages 1 through 7 compute the first microbatch's forward pass, and Stage 1 sits idle again while the pipeline drains at the end of the step. In a pipeline with N_{PP} stages and m microbatches, equation D.10 defines the **bubble fraction** as the proportion of total GPU-time spent idle:

$$b_{\text{pipe}} = \frac{N_{\text{PP}} - 1}{N_{\text{PP}} - 1 + m} \tag{D.10}$$

For balanced stages with negligible transfer overhead, this formula applies to both GPipe (where all forward passes complete before any backward pass begins) and 1F1B (one-forward-one-backward interleaving). The 1F1B schedule has the same asymptotic bubble fraction but uses less peak memory because it retires microbatches earlier.

The bubble fraction depends on the ratio of stages to microbatches, so adding more microbatches for a fixed number of stages amortizes the pipeline fill and drain across more useful work. As a rule of thumb, keeping bubble overhead at or below 20 percent requires $m \geq 4(N_{\text{PP}} - 1)$, while staying strictly below 20 percent requires $m > 4(N_{\text{PP}} - 1)$. The practical rule $m \geq 4N_{\text{PP}}$ is a convenient conservative target across typical stage depths and microbatch counts (table D.4).

Table D.4: Pipeline Bubble Fraction: Percentage of GPU-time lost to pipeline bubbles for various stage counts and microbatch counts. The conservative choice $m \geq 4N_{\text{PP}}$ keeps the modeled fraction below 20 percent.

Pipeline stages (N_{PP})	m **= 8**	m **= 16**	m **= 32**	m **= 64**
N_{PP} **= 4**	27.3%	15.8%	8.6%	4.5%
N_{PP} **= 8**	46.7%	30.4%	17.9%	9.9%

Deeper pipelines require proportionally more microbatches to maintain high accelerator utilization. When activation memory bounds limit the microbatch count, systems must adopt alternative scheduling strategies to mitigate bubble overhead.

D.4.2 Interleaved scheduling

When activation memory prevents a team from increasing m, splitting each GPU's layer assignment into smaller nonconsecutive model chunks lets the first microbatch reach the last device after fewer layers per stage. **Interleaved scheduling** (used in Megatron-LM) assigns V virtual stages to each physical GPU instead of one, so the pipeline has $N_{\text{PP}} \times V$ virtual stages while the bubble numerator retains the physical stage count N_{PP}, as equation D.11 shows:

$$b_{\text{interleaved}} = \frac{N_{\text{PP}} - 1}{N_{\text{PP}} - 1 + m \times V} \tag{D.11}$$

The microbatch term in the denominator is multiplied by V compared to equation D.10, so the bubble reduction approaches V when m dominates $N_{\text{PP}} - 1$. For $N_{\text{PP}} = 8$, $m = 16$, and $V = 4$, the bubble fraction drops from 30.4 percent to 9.9 percent, a 3.1× reduction that comes at the cost of more frequent, smaller communication between stages.

Interleaved scheduling highlights a recurring theme in distributed ML: communication granularity is a tunable knob. Smaller, more frequent transfers improve overlap and reduce idle time, but increase the aggregate latency overhead. The same trade-off governs gradient compression.

D.5 Compression Economics

> Systems Perspective D.5: Why this matters
>
> Gradient compression promises to reduce communication volume by 10–1000×, making AllReduce nearly free. Compression, however, has overhead (encoding and decoding time), and it only helps if that overhead is smaller than the communication time it saves. A quick break-even analysis prevents wasted effort on compression schemes that will not actually speed up training.

D.5.1 The compression equation

Compressing a 1 GB gradient tensor down to 16 MB cuts network transfer time on a 256-GPU ring from 40 ms to less than 2 ms, but only if the GPU kernels that quantize or sparsify the tensor finish in less than the 38 ms saved. Gradient compression replaces the original M-byte gradient with a compressed representation of size M/r_{comp}, where r_{comp} is the compression ratio. Equation D.12 expresses the total time with compression:

$$T_{\text{compressed}} = T_{\text{encode}} + T_{\text{transfer}}\left(\frac{M}{r_{\text{comp}}}\right) + T_{\text{decode}} \tag{D.12}$$

Compression is beneficial when $T_{\text{compressed}} < T_{\text{transfer}}(M)$ and the reduced precision or sparsity does not slow model convergence. Rearranging this inequality yields the break-even condition in equation D.13:

$$T_{\text{encode}} + T_{\text{decode}} < T_{\text{transfer}}(M) - T_{\text{transfer}}\left(\frac{M}{r_{\text{comp}}}\right) \tag{D.13}$$

In the bandwidth-dominated regime ($M \gg n^*$), the communication time scales linearly with message size, and the right-hand side simplifies to approximately $T_{\text{transfer}}(M) \times (1 - 1/r_{\text{comp}})$. For high compression ratios ($r_{\text{comp}} \gg 1$), this approaches $T_{\text{transfer}}(M)$, so compression is worthwhile as long as the codec overhead is less than the uncompressed communication time.

Napkin Math D.3: Worked example: Is compression worth it?

Setup: Ring AllReduce of a 1 GB gradient across 256 H100 GPUs on InfiniBand NDR, taking 40.6 ms uncompressed. Assume a collective-compatible compressed representation achieves 64× compression with 2 ms total encode + decode overhead.

Question: Does compression reduce the total AllReduce time?

Math:

- Uncompressed AllReduce: 40.6 ms
- Compressed communication: 1.4 ms
- Total with compression: 3.4 ms
- Speedup: 12×

Result: Compression shortens communication time by 12×. The 2 ms codec overhead is a small fraction of the 40.6 ms uncompressed baseline.

Systems insight: Compression saves communication time under the assumed payload ratio, codec cost, and unchanged collective structure. Independent top-k selections do not generally preserve that structure: sparse indices, differing supports, and merge operations must also be modeled. Lossy compression can require more iterations to reach the same accuracy, negating a per-iteration gain.

Break-even codec overhead limits (table D.5) scale with compression ratio and uncompressed AllReduce time. Higher compression ratios tolerate larger codec overheads, up to the bandwidth-dependent portion of the uncompressed collective. These break-even limits show that gradient compression delivers the highest net benefit on lower-bandwidth interconnects where transfer times dominate. On high-speed fabrics, the tolerable codec overhead budget is tightly constrained and requires lightweight quantization or sparsification algorithms.

Table D.5: Compression Break-Even Thresholds: Maximum tolerable codec overhead (encode + decode) for compression to reduce total time, computed as the actual collective-time difference $T_{\text{transfer}}(M) - T_{\text{transfer}}(M/r_{\text{comp}})$. The shared startup latency cancels, so even at high compression ratios it is not part of the codec budget.

Compression ratio (r_{comp})	AllReduce = 10 ms (max codec overhead)	AllReduce = 40 ms (max codec overhead)	AllReduce = 100 ms (max codec overhead)
4×	6.9	29.4	74.4
16×	8.7	36.8	93.0
64×	9.1	38.6	97.7
256×	9.2	39.1	98.8

Summary

The communication models in this appendix rest on two hardware parameters, the startup latency α and the bandwidth β of the α-β model (principle 5). They also rest on two workload parameters, the payload (n or M) and, for collectives, the participant count N. The crossover size $n^* = \alpha\beta$ separates latency-dominated from bandwidth-dominated transfers. Ring collectives carry a $(N-1)/N$ bandwidth factor and a latency term linear in N, while tree algorithms trade bandwidth for logarithmic latency, so message size relative to the crossover $M_{\text{crossover}}(N)$ selects the algorithm. Pipeline parallelism adds the bubble fraction $b_{\text{pipe}} = (N_{\text{PP}} - 1)/(N_{\text{PP}} - 1 + m)$, which the rule $m \geq 4N_{\text{PP}}$ keeps below 20 percent. Gradient compression pays only when codec overhead is smaller than

the communication time it saves and convergence is unaffected. Measuring α and β for a cluster therefore yields a first-order prediction of its communication overhead.

E

Reliability Foundations

Individual accelerators fail rarely, but scaling to thousands of nodes multiplies component failure rates until interruption becomes a continuous operating condition. At cluster scale, mean time between failures (MTBF) drops from years to hours, turning failure recovery into a core execution path. This appendix collects the quantitative reliability calculations behind the book's fault-tolerance chapters, deriving failure probabilities, Young-Daly optimal checkpoint intervals, and availability stacking models. In C^3 terms, it develops the coordination axis, where checkpointing, recovery, and availability determine how much of the fleet's compute and communication becomes useful work.

How to Use This Appendix

This appendix serves as a reference for moving from intuition ("failures happen more often at scale") to quantitative engineering decisions ("how often should one checkpoint?" or "how many spare nodes are needed?"). It translates component-level reliability metrics into cluster-wide checkpoint cadences and availability bounds.

The worked examples in this appendix use a 10,000-GPU training fleet, a larger variation on the 1,024-GPU Archetype A reference cluster, unless noted otherwise. The questions below map common operational problems to the sections that develop their analytical models:

- **"How often will something fail?"** Start with section E.1 and the MTBF cascade in section E.1.3.
- **"Where does the $1/N$ come from?"** The probability tools in section E.1.1 derive it from independent exponential lifetimes.
- **"How often to checkpoint?"** Use the Young-Daly model in section E.2.1 and the worked example in section E.2.3.
- **"How much time is lost to recovery?"** See the recovery anatomy in section E.3.1 and the runtime goodput loss in section E.3.2.
- **"When to use redundancy vs. checkpointing?"** Compare strategies in section E.4.1 and availability stacking in section E.4.2.
- **"What if checkpoints are not cheap relative to MTBF?"** Use the higher-order interval in section E.5.
- **"How can one flipped bit derail training?"** See the silent data corruption model in section E.6.

E.1 Failure Probability at Scale

Consider a training run on a 10,000-GPU cluster running for 30 days. The probability that at least one modeled node component fails during that time is effectively 100 percent. That answer determines whether the system needs fault tolerance as a core design requirement or as an optional feature.

Individual hardware components are reliable. A data center-grade GPU operates for tens of thousands of hours before failing. The physics of large-scale systems, however, works against reliability: every additional component is another opportunity for failure, and the aggregate failure rate scales linearly with component count. Component-level failure rates aggregate systematically into cluster-wide failure predictions.

E.1.1 Probability tools for failure analysis

A 10,000-GPU training job requires every accelerator and network card to stay healthy at the same time. At a 99.99 percent per-device daily survival rate, the job still has less than a 40 percent chance of finishing a single day uninterrupted. Four results from elementary probability carry that failure arithmetic throughout this book. Each assumes that components fail independently; correlated failures, such as a shared power feed or a common software bug, violate that assumption and must be modeled separately.

The first is the product rule for independent events. If a job needs all N of its components and each is available with probability $A_{\text{component}}$, the job's components are all available with probability

$$A_{\text{all}} = A_{\text{component}}^{N} \approx e^{-N(1-A_{\text{component}})} \quad \text{for } A_{\text{component}} \text{ close to } 1 \tag{E.1}$$

Equation E.1 is the **reliability collapse** behind the reliability gap in section 1.4.1. A per-component availability that looks excellent in isolation decays exponentially in N.

The second is its complement, the probability that at least one of N independent events occurs when each has probability p:

$$\Pr(\text{at least one}) = 1-(1-p)^{N} \approx Np \quad \text{when } Np \ll 1 \tag{E.2}$$

The linear approximation in equation E.2 holds only while Np is small; as N grows, the probability approaches 1. This arithmetic is why one slow worker paces a large synchronous job. A delay that rarely strikes any one worker strikes some worker of a large synchronous job on almost every step.

The third is the exponential lifetime model. A component with constant failure rate λ survives to time t with probability $e^{-\lambda t}$ and has mean time to failure $1/\lambda$. The model is memoryless. A component that has run for a year is no more likely to fail in the next hour than a new one, which is why it describes the useful-life phase and not burn-in or wear-out. The earliest of N independent exponential failures is itself exponential with rate $N\lambda$, so the time to the first failure among N identical components has mean $1/(N\lambda)$. That single fact is the origin of the $1/N$ scaling of system MTBF in equation E.5.

The fourth is the Poisson count. Under the same model, the number of failures in an interval T is Poisson distributed with mean T/MTBF, so the probability of no failure is $e^{-T/\text{MTBF}}$. Equation E.6 below is the complement of that event, and the expected number of failures per day in table E.2 is the Poisson mean.

E.1.2 Component failure rates

A single data center GPU fails roughly once every 5.7 years of continuous operation, a rate of about 0.00002 failures per hour that is too small to write conveniently in engineering tables. Reliability engineers therefore characterize components using two complementary metrics. The **failure in time (FIT)** rate counts failures per 10^9 device-hours of operation, a unit chosen because individual components fail so rarely that failures-per-hour would produce inconveniently small numbers. The reciprocal quantity, **mean time to failure (MTTF)**, calculates average lifetime in hours (equation E.3):

$$\text{MTTF} = \frac{10^9}{\text{FIT}} \tag{E.3}$$

Reference FIT rates and MTTF values (table E.1) characterize component reliability in typical GPU training nodes, grounded in large-fleet and warehouse-scale experience (Kokolis et al. 2025; Zu et al. 2024; Barroso et al. 2019). These values assume the steady-state "useful life" phase of the bathtub curve, where the failure rate is approximately constant, dominated neither by infant mortality (early life) nor by wear-out (end of life) (Klutke et al. 2003).

Table E.1: Component Failure Rates: Order-of-magnitude reference FIT/MTTF in the steady-state useful-life phase. Informed by large-GPU research-cluster analysis (Kokolis et al. 2025), TPUv4 supercomputer resiliency and operations (Zu et al. 2024), and warehouse-scale machine design (Barroso et al. 2019).

Component	FIT rate	MTTF	MTTF (years)	Typical failure mode
GPU	20,000	50,000 hours	5.7 years	Die defect, thermal fatigue
HBM	5,000	200,000 hours	22.8 years	Bit-flip accumulation, TSV
NIC	6,666.7	150,000 hours	17.1 years	Transceiver degradation
PSU	10,000	100,000 hours	11.4 years	Capacitor aging
PCIe Switch	5,000	200,000 hours	22.8 years	Solder joint, ESD damage
Optical Cable	20,000	50,000 hours	5.7 years	Fiber bend, connector wear
ToR Switch	3,333.3	300,000 hours	34.2 years	ASIC failure, fan bearing

Each component in isolation appears highly reliable; a GPU lasts 5.7 years on average. The difficulty begins with how a node behaves when many such components operate simultaneously.

E.1.3 The MTBF cascade

If one GPU or one network card inside an eight-GPU server fails during a synchronous collective, the entire server drops out of the training step just as surely as if its motherboard had lost power. The following model therefore treats a compute node as a series system, in which any modeled component failure interrupts it. A component enters the sum only when its failure alone stops the node, so redundant power supplies, spare links, and recoverable device errors stay out of it. For independent components with constant failure rates, equation E.4 sums the individual rates into the node-level failure rate:

$$\frac{1}{\text{MTBF}_{\text{node}}} = \frac{G_{\text{node}}}{\text{MTTF}_{\text{GPU}}} + \frac{n_{\text{nic}}}{\text{MTTF}_{\text{nic}}} + \frac{n_{\text{psu}}}{\text{MTTF}_{\text{psu}}} + \cdots \tag{E.4}$$

Think of each component as a ticking clock counting down to failure. A DGX H100 node with 8 GPUs and 8 network interface cards (NICs) has 16 independent clocks, and the node fails when the first clock reaches zero. More clocks mean a shorter expected wait. The node's 6 power supply units (PSUs) add no clocks. They run as a 4+2 redundant set, so the node keeps full performance with 2 of them failed, and a failed unit is normally replaced long before enough others fail to stop the node.

For a cluster of N_{nodes} identical nodes, the same series-system logic applies one level up. As equation E.5 shows, cluster MTBF divides the single-node MTBF by the number of nodes:

$$\text{MTBF}_{\text{system}} = \frac{\text{MTBF}_{\text{node}}}{N_{\text{nodes}}} \tag{E.5}$$

Reliability degrades linearly with component count at each level, and the levels compound, which is the **MTBF cascade**. A node with 4,687.5 hours MTBF sounds reliable. A cluster of 1,250 such nodes has an MTBF of 3.75 hours, so a failure every few hours is the expected steady state. Table E.2 carries that arithmetic across fleet sizes, with aggregate MTBF falling in inverse proportion to GPU count.

Table E.2: Cluster MTBF by Scale: As cluster size grows, the aggregate MTBF shrinks proportionally. At 10,000 GPUs, failures occur every few hours; at 100,000 GPUs, they occur continuously. Node configuration (DGX H100): 8 GPUs and 8 NICs per node in the series sum; its 6 PSUs are 4+2 redundant and add no term.

Cluster GPUs	Nodes	Cluster MTBF	Expected failures/day
256	32	146.5 hours	0.2
1,024	128	36.6 hours	0.7
2,048	256	18.3 hours	1.3
8,192	1,024	4.6 hours	5.2
10,000	1,250	3.8 hours	6.4
100,000	12,500	22.5 minutes	64

The transition from hundreds of GPUs to tens of thousands is a qualitative regime shift. About 6.1 days pass between failures at 256 GPUs. A 10,000-GPU cluster, in contrast, expects multiple failures per shift. At 100,000 GPUs, failures become a continuous background condition in which the system is never fully healthy.

E.1.4 Probability of failure during a job

A cluster MTBF of a few hours tells an operator how often nodes drop out on average, but a training team scheduling a 7-day or 30-day run needs the probability that at least one failure will interrupt that fixed window. Under the exponential failure model (constant failure rate), equation E.6 gives the probability of at least one failure during a job of duration T_{job} as the complement of the Poisson zero-failure probability over that window:

$$\Pr(\geq 1 \text{ failure}) = 1 - e^{-T_{\text{job}}/\text{MTBF}} \tag{E.6}$$

When $T_{\text{job}} \gg \text{MTBF}$, this probability approaches 1 rapidly, and it does so sooner on a larger cluster, whose MTBF is shorter (table E.3).

Table E.3: Probability of at Least One Failure: For large clusters and multi-day jobs, failure is a near-certainty. Any system operating in the bottom-right region of this table must treat fault tolerance as a core design requirement, not an optimization.

Cluster GPUs	1 day (24 h)	1 week (168 h)	30 days (720 h)
256	15.1%	68.2%	99.3%
1,024	48.1%	99.0%	> 99.9%
2,048	73.0%	> 99.9%	> 99.9%
8,192	99.5%	> 99.9%	> 99.9%
10,000	99.8%	> 99.9%	> 99.9%
100,000	> 99.9%	> 99.9%	> 99.9%

For any cluster above a few thousand GPUs running jobs longer than a day, the probability of experiencing at least one failure is effectively 100 percent. This is why chapter 7 treats fault tolerance as an architectural requirement rather than a defensive measure.

The exponential failure model assumes a constant failure rate, which holds during the steady-state useful-life phase. During burn-in (first few hundred hours) and wear-out (approaching end-of-life), failure rates are higher. In practice, fleet operators observe that newly deployed nodes show 2–3$\times$ higher failure rates in their first week, making burn-in testing essential before admitting nodes to production clusters. Because failures will interrupt long training jobs, the practical question is how to minimize the lost progress.

E.2 Checkpoint Optimization

On a 10,000-GPU cluster that fails every few hours, pausing all 10,000 accelerators to flush a multi-terabyte state snapshot to shared storage stalls training, yet spacing snapshots hours apart sacrifices thousands of GPU-hours of progress whenever a node drops out. Choosing a checkpoint cadence therefore trades guaranteed write overhead against expected rework after a crash (Young 1974; Daly 2006).

E.2.1 The Young-Daly model

During a synchronous checkpoint, writing state to storage takes time T_{write} (the **checkpoint cost**), during which no useful training occurs. Extending the interval between checkpoints amortizes that pause, but a failure that arrives at a random point inside the interval destroys on average half the interval's work. Equation E.7 states the **Young-Daly formula** that minimizes the sum of those two expected losses:

$$T^*_{\text{ckpt}} = \sqrt{2 \times T_{\text{write}} \times \text{MTBF}_{\text{system}}} \tag{E.7}$$

where T_{write} is the checkpoint write time in seconds and $\text{MTBF}_{\text{system}}$ is the cluster or job-level system MTBF in seconds.

The intuition behind equation E.7 is geometric-mean-like. When checkpoints are cheap relative to the MTBF ($T_{\text{write}} \ll \text{MTBF}_{\text{system}}$, the common case), the optimal interval sits between the two time scales. If checkpoints took zero time, the optimal cadence would be every step. If the system never failed, no checkpoints would be required. The square root interpolates between those extremes.

The formula assumes that failures follow an exponential distribution (memoryless property) and that checkpoint cost T_{write} is small compared to $\text{MTBF}_{\text{system}}$. Correlated outages, changing hazard rates, long writes, and asynchronous persistence can require a more detailed model. Under the first-order assumptions, evaluating equation E.7 for a concrete model requires only the cluster MTBF and the write duration T_{write}, which is set by the serialized state volume and storage bandwidth.

E.2.2 Checkpoint sizing

Resuming a mixed-precision training run after a node failure requires restoring not only the low-precision weights used in the forward pass, but also the full-precision master weights and optimizer moments accumulated across prior steps, which multiplies the serialized footprint and sets the write time T_{write} in equation E.7. For common mixed-precision training checkpoints with the adaptive moment estimation (Adam) optimizer, each parameter requires approximately 14 bytes of persistent state:

- 2 bytes for BF16 model weights
- 4 bytes for FP32 master weights
- 4 bytes for FP32 first moment (Adam m)
- 4 bytes for FP32 second moment (Adam v)

Equation E.8 aggregates these per-parameter costs into the total checkpoint footprint. The resulting byte volume scales linearly with model parameter count:

$$\text{Checkpoint Size} = P \times 14 \text{ bytes/parameter} \tag{E.8}$$

Here, P is the number of trainable model parameters whose persistent optimizer and weight state must be serialized. Transient activation buffers and unsharded temporary workspaces are excluded because they are recomputed after restore.

Mixed-precision Adam checkpoint sizing (table E.4) demonstrates how state scales from manageable gigabytes to frontier-scale terabytes. Even with high-bandwidth parallel storage, serializing these footprints requires tens of seconds to several minutes per snapshot:

Table E.4: Checkpoint Sizes for Mixed-Precision Adam Training: Each parameter requires about 14 bytes of persistent checkpoint state (model weights, FP32 master weights, and optimizer moments); gradients are normally transient and recomputed after restore rather than serialized as durable checkpoint state. Write times assume 100 GB/s aggregate storage bandwidth.

Model size	Checkpoint size	Write time at 100 GB/s
7B	98 GB	0.98 s
13B	182 GB	2 s
70B	980 GB	10 s
175B	2,450 GB	24 s
1T	14,000 GB	140 s

At frontier scale (175B+ parameters), these sizes reach the terabyte range. This makes checkpoint write time a significant cost that directly affects the Young-Daly optimal interval. The checkpoint strategies in chapter 7 discuss techniques for reducing T_{write} (asynchronous checkpointing, incremental deltas, and distributed storage), all of which improve the Young-Daly result by shrinking the T_{write} factor under the square root.

E.2.3 Calculating the optimal checkpoint interval

Combining the checkpoint write time with the 10,000-GPU MTBF gives a concrete checkpoint cadence. The worked example below walks through the two-step calculation and shows how doubling the cluster size shifts the optimum.

Napkin Math E.1: Young-Daly: 175B model on a 10,000-GPU cluster

Setup: Consider training a 175B-parameter model on a 10,000-GPU cluster. The cluster MTBF is 3.75 hours (table E.2). The checkpoint size is 2,450 GB, and the parallel storage system writes at 100 GB/s.

Step 1: Checkpoint write time (T_{write}).

$$T_{\text{write}} = \frac{\text{Checkpoint Size}}{\text{Write Bandwidth}} = \frac{2,450 \text{ GB}}{100 \text{ GB/s}} = 24.5 \text{ s}$$

Step 2: Apply the Young-Daly formula (equation E.7).

$$T^*_{\text{ckpt}} = \sqrt{2 \times T_{\text{write}} \times \text{MTBF}_{\text{system}}} = \sqrt{2 \times 24.5 \text{ s} \times 3.75 \text{ h} \times 3,600 \text{ s/h}} = 13.6 \text{ min}$$

Interpretation: The optimal checkpoint interval is approximately 14 minutes. The overhead from checkpointing alone is $T_{\text{write}}/T^*_{\text{ckpt}} \approx 3.0$ percent of training time.

Systems insight: If the cluster were doubled to 20,000 GPUs, the MTBF would halve and the optimal interval would shrink to 9.6 minutes. The system therefore checkpoints more frequently because failures happen more often. Larger clusters are faster but must interrupt themselves more often to protect progress, which is the tension at scale.

When T_{write} approaches $\text{MTBF}_{\text{system}}$ (checkpoint cost approaches MTBF), checkpoint and rework overheads become so large that checkpoint/restart alone may fail to maintain useful forward progress, requiring the alternative strategies in section E.4.1. Even when $T_{\text{write}} \ll \text{MTBF}_{\text{system}}$, the Young-Daly cadence prices only the periodic write pause and the average step replay after a fault; a real cluster outage also stalls on failure detection, node replacement, and checkpoint reload before replay can begin.

E.3 Recovery Budgets

When a node drops out of a 10,000-GPU training step, every surviving accelerator idles while the control plane times out the dead heartbeat, allocates a replacement node, streams terabytes of checkpoint state back into GPU memory, and recomputes the steps lost since the last snapshot. Pricing how much wall-clock capacity those interruptions consume requires decomposing the outage into its constituent phases and translating the total stall time into runtime goodput.

E.3.1 The anatomy of recovery time

A single hardware fault triggers four sequential recovery phases between the instant a worker stalls and the moment the job returns to new training steps. Equation E.9 sums them into total recovery time:

$$T_{\text{recovery}} = T_{\text{detect}} + T_{\text{restart}} + T_{\text{load}} + T_{\text{replay}} \tag{E.9}$$

The terms are all durations: failure detection, replacement scheduling, checkpoint reload, and replay of lost training work since the last checkpoint. Their sum is the wall-clock time before the job returns to productive training.

Table E.5 breaks recovery into the phases that determine how long the cluster remains below full productivity after a failure. Because detection, rescheduling, checkpoint reload, and replay appear as separate rows, the table shows which operational bottlenecks dominate total downtime.

Table E.5: Recovery Time Breakdown: Each phase contributes to the total time between failure and full-speed resumption. For the 10,000-GPU, 175B-model scenario, replay dominates because it recomputes work lost since the last checkpoint.

Phase	Typical duration	What happens
T_{detect}	30 s	Heartbeat timeout expires; worker declared failed (confirmation adds to the budget in chapter 7)
T_{restart}	60 s	Replacement node allocated from spare pool
T_{load}	24.5 s	Checkpoint read from storage into GPU memory
T_{replay}	~6.8 min	Recompute training steps since last checkpoint
Total T_{recovery}	~9 min	System fully productive again

In this sequence, replay time T_{replay} typically dominates and is controlled directly by the checkpoint interval, since on average half the interval must be replayed. This creates a reinforcing loop with the Young-Daly formula. Shorter intervals mean less replay but more checkpoint overhead, and the formula finds the minimum of this sum. The other phases offer distinct engineering levers. More aggressive heartbeat intervals shorten T_{detect}, at the cost of false positives. Preallocated hot spare nodes shorten T_{restart} through the fleet orchestration policies covered in chapter 8. Higher storage bandwidth or smaller checkpoints shorten T_{load}, which motivates the compression and sharding techniques discussed in chapter 7.

E.3.2 Runtime goodput under failures

Every minute a cluster spends writing a periodic snapshot or executing the four-phase recovery pipeline of table E.5 is billed wall-clock allocation that advances the final model by zero training steps. Comparing the **committed step rate** (the steps a checkpoint retains for the final model) against the **fault-free step rate** (the rate the configured cluster could sustain without checkpoint or recovery interruptions) defines the runtime goodput of a job whose only losses are checkpoints and failures:

$$\eta_{\text{runtime}} \approx \frac{\text{Committed Useful-Step Rate}}{\text{Fault-Free Step Rate}} \approx \frac{T_{\text{useful}}}{T_{\text{wall}}}$$

Section 9.5 defines runtime goodput as one factor of ML goodput and also charges startup and input stalls to it, which this appendix does not model. The modeled gap between the two rates comes from two sources:

1. **Checkpoint overhead** ($\sim$ 3.0 percent): Training pauses during each checkpoint write.

2. **Recovery and rework overhead** ($\sim$ 3.9 percent): Time lost to detection, rescheduling, reloading, and replaying work discarded after each failure.

At 10,000-GPU scale, failures and checkpointing together can consume a substantial fraction of wall time. If they consume a fraction f_{overhead}, a fixed amount of useful work takes $1/(1 - f_{\text{overhead}})$ times its fault-free duration, not merely that duration plus f_{overhead} of it. The fleet orchestration strategies in chapter 8 and chapter 12 focus on narrowing this gap, because every percentage point of overhead recovered saves dollars and shortens training time.

> Systems Perspective E.1: The hidden cost of scale
>
> A common misconception is that doubling cluster size halves training time. In practice, doubling from 5,000 to 10,000 GPUs halves the MTBF and increases failure-related overhead. The effective speedup is less than 2×, and at extreme scale, adding more GPUs can actually increase wall-clock time if the fault tolerance mechanisms cannot keep pace. The serial overhead of recovery bounds the benefit of parallelism, a reliability analogue of Amdahl's Law.

E.4 Strategy Selection

Pausing a 10,000-GPU batch training job for several minutes to reload a checkpoint and replay lost steps is an acceptable throughput tax, whereas stalling a live inference service for even a few seconds violates user-facing latency contracts and drops production traffic. Because batch training and online serving tolerate interruption on completely different time scales, fleet architectures choose among checkpoint/restart, active redundancy, and elastic resizing according to the recovery latency, spare capacity, and state-management complexity each workload can afford.

E.4.1 Checkpoint/restart vs. redundancy vs. elastic training

Whether a fault-tolerance mechanism can pause the entire job to replay work, must fail over within milliseconds, or must continue on a shrunk allocation when replacement nodes are unavailable separates the three primary recovery architectures. **Checkpoint/restart** periodically saves full system state and rolls back to the last checkpoint after failure. It is the workhorse of large-scale training: conceptually simple, well-understood, and effective when MTBF is much larger than checkpoint cost. The weakness is that recovery requires stopping all workers and replaying lost computation.

Redundancy maintains alternate copies of state or computation. A surviving serving replica can take over only when routing, readiness, capacity, and the required request state meet the latency budget. Dedicated duplicate or triplicate execution costs 2–3× the compute resources; active-active serving instead shares normal traffic across replicas while reserving failover headroom. The justified overhead depends on the workload and recovery contract.

Elastic training allows a job to resume with fewer workers when the smaller allocation can hold its state and support its parallel decomposition. Membership changes can still require a coordinated restart and checkpoint reload; TorchElastic restarts the worker group. Workers may be added back when capacity becomes available. The recovery path must restore missing shards, preserve data coverage, and either preserve the effective batch size or validate a changed optimization policy. Elasticity can reduce waiting for replacement capacity, but does not guarantee a shorter interruption. The three strategies trade recovery latency against resource overhead and implementation complexity, and each fits a different workload (table E.6).

Table E.6: Fault Tolerance Strategy Comparison: The latency and overhead ranges are illustrative, not guarantees; elasticity does not by itself require spare capacity. Real-world systems often combine strategies: checkpoint/restart for training with redundancy for the metadata service and checkpoint storage layer.

Criterion	Checkpoint/restart	Redundancy	Elastic training
Recovery latency	Minutes (replay)	Milliseconds (failover)	Restart and resize dependent
Resource overhead	~3–13% (storage + IO)	100–200% (replicas)	Resize and standby policy dependent

Criterion	Checkpoint/restart	Redundancy	Elastic training
Workload fit	Training (batch)	Serving (online)	Training (long-running)
Implementation	Simple	Moderate	Complex
State management	Periodic snapshots	Continuous replication	Distributed with resharding
Failure mode	Job pauses, replays	Transparent to user	Pause, restore, then resume

E.4.2 The availability stacking formula

A single serving replica with 99 percent availability ("two nines") still accumulates more than 87 hours of downtime per year, far exceeding the outage budget of production inference endpoints that target 99.9 percent ("three nines") or 99.99 percent ("four nines") uptime. Redundancy closes that gap by running k independent replicas so that surviving instances absorb traffic when one replica fails. Under an idealized model with reliable routing and enough capacity on any surviving replica, the system is unavailable only when all replicas are simultaneously down. Equation E.10 captures this, with A_{single} denoting the availability of a single replica and k the number of replicas:

$$A_{\text{system}} = 1 - (1 - A_{\text{single}})^k \tag{E.10}$$

Evaluating availability stacking with independent replicas (table E.7) quantifies the gain from adding serving capacity. Each additional independent replica multiplies the outage probabilities together, cutting expected annual downtime by two orders of magnitude per replica at a 99 percent baseline.

Table E.7: Availability Stacking with Independent Replicas: Starting from a single-replica availability of 99 percent, each additional replica dramatically reduces expected downtime. Assumes replica failures are independent.

Replicas k	System availability	Nines	Downtime per year
1	99.00%	2.0	87.6 hours
2	99.9900%	4.0	52.6 minutes
3	99.9999%	6.0	31.5 seconds

Stacking independent redundant replicas achieves large availability gains. Two replicas that are each available 99 percent of the time yield 99.99 percent availability, reducing annual downtime from 87.6 hours to 52.6 minutes. This is why inference serving systems almost universally deploy multiple replicas behind a load balancer. The cost of an extra replica is small compared to the business value of four-nines availability.

The result, however, depends on the independence assumption. Correlated failures, such as power outages affecting an entire rack, software bugs triggered by a specific input, or network partitions isolating a failure domain, defeat availability stacking. This is why chapter 7 emphasizes failure domain isolation: replicas must be placed in different racks, different power zones, and ideally different data centers to reduce shared failure modes. Placement alone does not eliminate common software, control-plane, or regional dependencies.

E.5 Checkpoint Intervals Beyond the First-Order Approximation

When a multi-terabyte checkpoint takes several minutes to flush to shared storage and a 100,000-GPU cluster fails every half hour, the checkpoint write is no longer negligible next to the failure interval. The cluster can then suffer a second crash while still writing the snapshot. Young's first-order approximation is most accurate when checkpoint time is small relative to the failure interval; Daly's analysis develops higher-order checkpoint models for regimes where write time is a substantial fraction of MTBF (Daly 2006). The simplified renewal model below shows why finite checkpoint time changes the optimum without claiming to reproduce Daly's full recovery model.

E.5.1 Renewal model and small-cost expansion

Suppose a worker fails halfway through flushing a 60-second checkpoint snapshot; the job loses both the training steps computed since the previous snapshot and the 30 seconds spent writing the incomplete checkpoint. Let T_{ckpt} be useful compute time between completed checkpoints. Assume Poisson failures with rate $1/\text{MTBF}_{\text{system}}$, identical failure exposure during computation and writing, and zero detection, restart, or reload delay. A failed attempt loses its partial progress and repeats the interval. The expected elapsed time to complete one uninterrupted compute-plus-checkpoint attempt is:

$$\mathbb{E}[T_{\text{total}}] = \text{MTBF}_{\text{system}} \left(e^{\frac{T_{\text{ckpt}}+T_{\text{write}}}{\text{MTBF}_{\text{system}}}} - 1 \right) \tag{E.11}$$

Writing $T_{\text{att}} = T_{\text{ckpt}} + T_{\text{write}}$ for the attempt duration, the renewal equation is $\mathbb{E}[T_{\text{total}}] = \mathbb{E}[\min(F, T_{\text{att}})] + \Pr(F < T_{\text{att}})\mathbb{E}[T_{\text{total}}]$, where F is an exponential failure time. Since $\mathbb{E}[\min(F, T_{\text{att}})] = \text{MTBF}_{\text{system}}(1 - e^{-T_{\text{att}}/\text{MTBF}_{\text{system}}})$, solving the equation gives equation E.11. Expanding the exponential through its quadratic term then gives the wasted time per completed interval:

$$\mathbb{E}[T_{\text{total}}] - T_{\text{ckpt}} \approx T_{\text{write}} + \frac{(T_{\text{ckpt}} + T_{\text{write}})^2}{2 \cdot \text{MTBF}_{\text{system}}} \tag{E.12}$$

When $T_{\text{write}}/\text{MTBF}_{\text{system}}$ is small, dividing equation E.12 by T_{ckpt} leaves the leading terms $T_{\text{write}}/T_{\text{ckpt}} + T_{\text{ckpt}}/(2 \cdot \text{MTBF}_{\text{system}})$ of the wasted-time fraction. Minimizing those terms recovers Young's interval $\sqrt{2 \cdot T_{\text{write}} \cdot \text{MTBF}_{\text{system}}}$. For finite checkpoint costs, differentiate the unexpanded elapsed time per unit useful work, $\mathbb{E}[T_{\text{total}}]/T_{\text{ckpt}}$, instead. Its stationary point satisfies:

$$e^{(T^*_{\text{ckpt}}+T_{\text{write}})/\text{MTBF}_{\text{system}}} \left(1 - \frac{T^*_{\text{ckpt}}}{\text{MTBF}_{\text{system}}} \right) = 1, \qquad 0 < T^*_{\text{ckpt}} < \text{MTBF}_{\text{system}} \tag{E.13}$$

For positive checkpoint cost, the equivalent derivative numerator

$$e^{(T_{\text{ckpt}}+T_{\text{write}})/\text{MTBF}_{\text{system}}}(T_{\text{ckpt}}/\text{MTBF}_{\text{system}} - 1) + 1$$

increases from a negative value at zero to a positive value at $\text{MTBF}_{\text{system}}$. A one-dimensional root search on equation E.13 therefore finds the unique optimum of this simplified model.

Choosing the appropriate approximation

The first-order formula is a useful planning estimate when writes are short compared with MTBF. When that condition fails, use a model that includes the deployment's checkpoint semantics, recovery delays, and failure exposure. The numerical condition in equation E.13 covers only the stated zero-restart-delay renewal model; asynchronous writes and correlated failures require different assumptions. All of these checkpoint and redundancy models assume that a hardware fault announces itself by stopping a worker or failing a health check, yet some of the most damaging faults at fleet scale corrupt tensor values silently without halting execution.

E.6 Silent Data Corruption in Attention

Not all hardware faults manifest as explicit system crashes, ECC double-bit errors, or kernel traps. **Silent data corruption (SDC)** occurs when single-bit errors or arithmetic logic unit (ALU) faults slip past parity checks, corrupting tensor values in SRAM, high-bandwidth memory (HBM), or register files without triggering an exception. In deep learning workloads, SDC is particularly dangerous because single corrupted floating-point values propagate through non-linear operations, destabilizing fleet-wide training (Dixit et al. 2021).

E.6.1 Exponent bit flip model in FP16

Consider an IEEE 754 half-precision floating-point (FP16) tensor element representing a query entry q_{im}, activation value, or weight matrix entry. FP16 allocates 16 total bits: 1 sign bit (s), 5 exponent bits (e), and 10 mantissa bits (m). For a finite normal value, with exponent field $1 \le e \le 30$, the scalar x represented by the bitstring (s, e, m) is given by:

$$x = (-1)^s \times 2^{e-15} \times \left(1 + \frac{m}{1024}\right)$$

Suppose a cosmic ray, voltage drop, or SRAM failure causes a single bit flip in the most significant exponent bit (bit 14). The exponent value changes from e to $e' = e \pm 16$:

$$x_{\text{corrupt}} = (-1)^s \times 2^{(e\pm16)-15} \times \left(1 + \frac{m}{1024}\right) = x \times 2^{\pm16}$$

If both the original and corrupted exponent fields encode finite normal values, the magnitude increases by a factor of $2^{16} = 65{,}536$ or decreases by its reciprocal. Exponent fields of zero or all ones instead invoke the subnormal, zero, infinity, or NaN encodings, so the displayed normal-value formula does not apply.

E.6.2 Softmax saturation and error propagation

When a single-bit exponent flip inflates one query entry by a factor of $65{,}536$, the dot products in that token's attention row can jump by orders of magnitude and collapse its softmax distribution onto a single key. To trace how that perturbation propagates through one attention head, let $\mathbf{S} = \mathbf{Q}\mathbf{K}^T/\sqrt{d_{\text{head}}} \in \mathbb{R}^{S\times S}$ be the pre-softmax score matrix for sequence length S and head dimension d_{head}. Suppose an SDC event corrupts one element of token row i of the query matrix $\mathbf{Q}$, and write the resulting perturbation as $\Delta\mathbf{Q}$, with only that element nonzero:

$$\mathbf{S}_{\text{corrupt}} = \frac{(\mathbf{Q} + \Delta\mathbf{Q})\mathbf{K}^T}{\sqrt{d_{\text{head}}}} = \mathbf{S} + \Delta\mathbf{S}$$

where $\Delta\mathbf{S}$ is nonzero only in row i, with entries $\Delta S_{ij} = \Delta\mathbf{q}_i \cdot \mathbf{k}_j/\sqrt{d_{\text{head}}}$ for the corrupted query row $\Delta\mathbf{q}_i$ and key row $\mathbf{k}_j$. Passing the perturbed score row through the normalization step shifts the entire row's attention weights. The attention probability perturbation for token row i is therefore:

$$\Delta\mathbf{a}_i = \text{softmax}\left(\mathbf{S}_{i,\cdot} + \Delta\mathbf{S}_{i,\cdot}\right) - \text{softmax}\left(\mathbf{S}_{i,\cdot}\right)$$

A large corrupted query component can strongly change several attention logits, with the signs and magnitudes determined by the corresponding key components. If one resulting logit exceeds every other logit by an increasingly large margin, the attention distribution approaches a one-hot vector:

$$\text{softmax}(\mathbf{S}_{i,\cdot} + \Delta\mathbf{S}_{i,\cdot})_j \to 1, \qquad \text{softmax}(\mathbf{S}_{i,\cdot} + \Delta\mathbf{S}_{i,\cdot})_{k\neq j} \to 0$$

This is a conditional saturation limit, not an inevitable consequence of a bit flip. The selected key depends on the corrupted logits; a corrupted query element does not identify one corrupted key. The softmax derivative is $\partial p_j/\partial S_{ik} = p_j(\delta_{jk} - p_k)$, which tends toward zero under saturation rather than diverging. Other paths through the computation can still produce large or nonfinite activations and gradients. If invalid state enters the optimizer or a collective, it can propagate across ranks. Detection and recovery must therefore verify state before resuming from a clean checkpoint.

Summary

Failure is a physical constraint of scale. Under independent, constant component failure rates, a single GPU's mean lifetime is 5.7 years. Yet the modeled 10,000-GPU cluster fails on average every 3.75 hours, because node failure rates sum across components and cluster MTBF divides by node count. The probability that a job is interrupted grows with its duration relative to that interval, so long runs on large clusters must plan for recovery rather than hope to avoid it.

Recovery has four modeled phases (detection, rescheduling, reloading, and replay), and replay is controlled by the checkpoint interval. The Young-Daly formula $T^{*}_{\text{ckpt}} = \sqrt{2T_{\text{write}}\,\text{MTBF}_{\text{system}}}$ (principle 13) balances write overhead against expected rework when writes are short compared with the MTBF. Section E.5 examines finite-cost effects under a simplified renewal model. Checkpoint/restart suits batch training, redundancy suits latency-sensitive serving, and elastic training sits between them. Independent replicas reduce unavailability multiplicatively, to $(1 - A_{\text{single}})^{k}$ for k replicas, but only to the extent that failure-domain isolation keeps their failures independent.

F

Inference Foundations

Serving a model at scale is governed by the mathematics of waiting lines and the physics of the key-value cache. This appendix develops the analytical serving models that the book's inference and operations chapters draw on, establishing how queuing dynamics and attention-cache memory arithmetic combine into a single serving constraint under strict latency service-level objectives (SLOs). In C^3 terms, inference turns compute capacity, communication with memory and clients, and coordination across queues into one serving constraint.

How to Use This Appendix

This appendix is a quantitative reference for serving systems. Use the queuing surrogate to screen batch sizes and estimate mean delay under its stated assumptions, then validate tail latency with measurements or simulation. Use the key-value cache arithmetic to size memory or to determine when the cache, rather than compute, caps throughput.

- **When a serving tier misses its latency target under load**: Use the queuing model in section F.1.1 to separate waiting, batch-formation, and service time.
- **When choosing a batch size**: Use the worked example in section F.1.1.4 and the decision framework in section F.1.1.5.
- **When memory, not compute, caps concurrency**: Use the KV-cache sizing in section F.1.2.

F.1 Serving Performance Analysis

When a GPT-3 serving replica receives a burst of 100 queries per second, an aggressive batch size can saturate high-bandwidth memory with key-value states while a conservative batch size causes the arrival queue to grow without bound. Serving performance analysis therefore has two coupled layers. The queue determines how requests wait, batch, and shed under stochastic arrivals; the key-value cache determines how much live state the accelerator can hold while meeting the latency budget. Because a batch cannot execute until it both clears the arrival queue and fits in device memory, sizing a serving replica requires solving the queuing delay curve first and then bounding candidate batch sizes against key-value cache capacity.

F.1.1 Queuing theory for batched inference

Processing requests one at a time leaves GPU tensor cores starved for work, yet waiting to accumulate a batch of 64 requests adds tens of milliseconds of collection delay before the first kernel even launches. Batching trades waiting time for throughput, and production systems require quantitative analysis to choose batch sizes under stochastic arrival patterns. Queuing theory provides the mathematical framework to derive these operating points and to explain why certain batch sizes outperform others under specific load conditions.

F.1.1.1 Queuing notation and the M/G/1 batch surrogate

The **M/G/c/K notation** names four characteristics that a production GPU-serving model may need to represent. Together, these four fields specify the arrival process, the service-time distribution, the server count, and the queue capacity:

- **M (Markov arrivals)**: Requests arrive according to a Poisson process with rate λ_{arr}. This models the memoryless property of user requests, where arrival of one request does not predict the timing of the next.

- **G (General service distribution)**: Service times follow a general distribution, not restricted to exponential. GPU inference times depend on batch size and show deterministic components (compute) mixed with stochastic variation (memory contention, kernel scheduling).
- **c (Number of servers)**: The system has c parallel GPU workers, each capable of serving requests independently.
- **K (System capacity)**: At most K requests may be present, including those in service. Requests arriving to a full system are rejected (load shedding).

The notation comes from the same systems-performance lineage that sized telephone networks before it was applied to GPU clusters. Its parameters map directly onto the arrival, execution, and admission controls of an inference server.

> Systems Perspective F.1: Queuing theory and performance
>
> Queuing theory, developed by Agner Krarup Erlang in 1909 for telephone network analysis, remains foundational to systems performance engineering (Erlang 1909). The same mathematical framework that sized telephone exchanges now informs GPU cluster capacity. Standard systems texts develop the component queuing models and approximations on which this notation builds (Jain 1991; Kleinrock 1975).

For a batched inference system with batch size B, service time becomes a function of batch size. Let $T_{\text{svc}}(B)$ denote the time to process a batch of B requests. From the physics of batching (section 10.6), equation F.1 models this as:

$$T_{\text{svc}}(B) = T_{\text{fixed}} + t_{\text{req}} \cdot B \tag{F.1}$$

where T_{fixed} represents fixed service overhead (kernel launch, weight loading from high-bandwidth memory (HBM)) and t_{req} represents marginal per-request computation time. This linear model captures the first-order behavior observed in production systems, though actual service times may show slight sublinearity due to memory bandwidth saturation at large batch sizes.

The **effective service rate** for batched processing divides the batch size by the batch execution time. That ratio determines the maximum sustained request throughput at a given batch size:

$$\mu_{\text{eff}}(B) = \frac{B}{T_{\text{svc}}(B)} = \frac{B}{T_{\text{fixed}} + t_{\text{req}} \cdot B}$$

The effective rate increases with batch size, approaching the asymptotic limit $1/t_{\text{req}}$ as $B \to \infty$. Once fixed launch and weight-loading overheads are amortized, marginal per-request work sets the throughput ceiling.

F.1.1.2 Response time analysis

Amortizing fixed launch and weight-loading overhead raises server throughput, but a client experiences end-to-end latency rather than kernel throughput alone. Before a request ever executes on the accelerator, it must wait for prior batches to clear and pause in the staging buffer while peer requests arrive to complete its own batch, splitting total response time t into three sequential components:

$$\mathbb{E}[t] = \mathbb{E}[W_q] + \mathbb{E}[L_{\text{lat,form}}] + \mathbb{E}[T_{\text{svc}}]$$

where:

- $\mathbb{E}[W_q]$ is the expected waiting time in queue before joining a batch
- $\mathbb{E}[L_{\text{lat,form}}]$ is the expected per-request batch-accumulation delay before dispatch
- $\mathbb{E}[T_{\text{svc}}]$ is the expected inference time once the batch executes

For dynamic batching with maximum wait time T_{window} and maximum batch size B_{max}, the first request triggers the batching window. Let $Q_{\text{arr}} \sim \text{Poisson}(\lambda_{\text{arr}} T_{\text{window}})$ count subsequent arrivals during that window and let $k = B_{\text{max}} - 1$. For $B_{\text{max}} > 1$, the realized batch size is $B = \min(B_{\text{max}}, 1 + Q_{\text{arr}})$, including the trigger request, so:

$$\mathbb{E}[B] = \sum_{n=0}^{k-1} (n+1)\Pr(Q_{\text{arr}} = n) + B_{\max}\Pr(Q_{\text{arr}} \geq k)$$

Let τ_k be the Erlang-distributed time of the kth subsequent arrival. A full batch closes when $\tau_k \leq T_{\text{window}}$, while a timeout batch closes at T_{window} with $Q_{\text{arr}} = n < k$. The mean per-request accumulation delay within each realized batch is:

$$\mathbb{E}[L_{\text{lat,form}} \mid \text{full}] = \frac{1}{2}\mathbb{E}[\tau_k \mid \tau_k \leq T_{\text{window}}]$$

$$\mathbb{E}[L_{\text{lat,form}} \mid Q_{\text{arr}} = n,\ \text{timeout}] = \frac{T_{\text{window}}}{2}\left(1 + \frac{1}{n+1}\right), \qquad 0 \leq n < k$$

Weighting each batch outcome by its realized request count gives the mean delay of a randomly selected request. This weighted average accounts for the fact that larger batches contain a proportionally larger share of the arriving queries:

$$\mathbb{E}[L_{\text{lat,form}}] = \frac{B_{\max}\Pr(Q_{\text{arr}} \geq k)\mathbb{E}[\tau_k/2 \mid \tau_k \leq T_{\text{window}}] + \sum_{n=0}^{k-1}(n+1)\Pr(Q_{\text{arr}} = n)\mathbb{E}[L_{\text{lat,form}} \mid Q_{\text{arr}} = n,\ \text{timeout}]}{\mathbb{E}[B]}$$

When batches reliably fill before timeout, this expression approaches $\mathbb{E}[\tau_k]/2 = (B_{\max} - 1)/(2\lambda_{\text{arr}})$. For $B_{\max} = 1$, the batch contains only its trigger request and accumulation delay is zero.

F.1.1.3 Optimal batch size derivation

Increasing the batch size B pulls the three latency components in opposite directions: a larger batch raises the effective service rate and drains the arrival queue faster, yet it forces every request to wait longer during batch formation and adds t_{req} per additional request during execution. The optimal batch size B minimizes expected response time across those competing terms subject to throughput stability:

$$B = \arg\min_B \mathbb{E}[t(B)] \quad \text{subject to} \quad \mu_{\text{eff}}(B) > \lambda_{\text{arr}}$$

The constraint ensures system stability so that the effective service rate strictly exceeds the arrival rate. Without that margin, the arrival queue grows without bound under stochastic traffic bursts.

For a high-traffic regime in which batches of size B reliably fill before timeout, substituting the response time components gives a closed-form objective. Each term in the sum responds differently as B increases:

$$\mathbb{E}[t(B)] = \mathbb{E}[W_q(B)] + \frac{B-1}{2\lambda_{\text{arr}}} + T_{\text{svc}}(B)$$

For a tractable waiting-time estimate, use an M/G/1 surrogate that treats batches as Poisson "super-requests." If the request arrival rate is λ_{arr} and the average batch size is B, this surrogate uses batch arrival rate $\lambda_{\text{batch}} = \lambda_{\text{arr}}/B$. This is not the exact process produced by grouping Poisson request arrivals, because fixed-size grouping yields Erlang inter-batch times; finite K or $c > 1$ requires a bulk-service model or simulation. Within the explicit surrogate, the Pollaczek-Khinchine formula gives:

$$\mathbb{E}[W_q] = \frac{\lambda_{\text{batch}} \cdot \mathbb{E}[T_{\text{svc}}^2]}{2(1 - \rho_{\text{serv}})}$$

where $\rho_{\text{serv}} = \lambda_{\text{batch}} \cdot \mathbb{E}[T_{\text{svc}}] = \lambda_{\text{arr}} \cdot \mathbb{E}[T_{\text{svc}}]/B$ is the server utilization. The second moment $\mathbb{E}[T_{\text{svc}}^2]$ captures service time variability.

For the linear service time model with deterministic service in this surrogate, the second moment simplifies to $T_{\text{svc}}(B)^2$. Substituting that deterministic service time into the Pollaczek-Khinchine relation gives the explicit waiting-time curve:

$$\mathbb{E}[W_q(B)] = \frac{\lambda_{\text{arr}} \cdot T_{\text{svc}}(B)^2}{2B(1 - \lambda_{\text{arr}} \cdot T_{\text{svc}}(B)/B)}$$

For production sizing, teams often first enforce a utilization headroom target rather than solving the full latency minimization problem. With $T_{\text{svc}}(B) = T_{\text{fixed}} + t_{\text{req}}B$, the constraint $\rho_{\text{serv}} = \lambda_{\text{arr}} T_{\text{svc}}(B)/B \leq \rho_{\text{target}}$ yields the target-utilization batch bound (equation F.2):

$$B_{\text{util}}(\rho_{\text{target}}) = \frac{\lambda_{\text{arr}} T_{\text{fixed}}}{\rho_{\text{target}} - \lambda_{\text{arr}} t_{\text{req}}} \tag{F.2}$$

where ρ_{target} is the target utilization (typically 0.7–0.8 for production systems to maintain latency headroom) and $\rho_{\text{target}} > \lambda_{\text{arr}} t_{\text{req}}$.

Systems Perspective F.2: Target-utilization batch sizing

The target-utilization batch bound (equation F.2) shows that the minimum utilization-feasible batch size grows with fixed overhead and arrival rate, through three dependencies:

1. **Higher traffic** (λ_{arr}): Required batch size increases.
2. **Higher fixed overhead** (T_{fixed}): Required batch size increases to amortize overhead.
3. **Higher target utilization** (ρ_{target}): Required batch size decreases, but with less latency headroom.

This utilization-constrained sizing explains why large language models (LLMs) (high T_{fixed} from weight loading) benefit from larger batches than vision models (low T_{fixed}), even at the same arrival rate.

F.1.1.4 GPT-3 serving at 100 QPS

Consider a hypothetical prefill-stage sizing exercise for a GPT-3 class model (175B parameters). The following assumed service-time coefficients illustrate the queuing calculation; they are not a measured GPT-3 benchmark:

- Arrival rate: $\lambda_{\text{arr}} = 100$ requests/second
- Hardware: 8× A100 GPUs (one DGX A100 node) with tensor parallelism
- Fixed prefill-stage service overhead: $T_{\text{fixed}} = 50$ ms
- Marginal prefill-stage service time: $t_{\text{req}} = 0.5$ ms per request
- Target utilization: $\rho_{\text{target}} = 0.75$

For the prefill phase (processing the input prompt), service time follows the linear batch model of equation F.1. Substituting the fixed 50 ms overhead and 0.5 ms marginal cost gives the per-batch execution time:

$$T_{\text{svc}}(B) = 50 + 0.5 \cdot B \text{ ms}$$

Applying equation F.2 yields the minimum batch size that meets the 75 percent target utilization while leaving 25 percent headroom. The numerator captures the fixed overhead rate and the denominator captures the available utilization margin:

$$B_{\text{util}} \approx \frac{100 \times 0.05}{0.75 - 100 \times 0.0005} = \frac{5}{0.70} \approx 7.14$$

Rounding up gives $B = 8$ as the smallest batch size that meets the target utilization headroom. Evaluating the surrogate response-time expression under the full-batch assumption then shows how larger batches reduce utilization and queuing delay but add batch-accumulation delay (table F.1).

Table F.1: Batch Size Impact on GPT-3 Serving Performance: Batch sizes of 5 or fewer cannot sustain 100 QPS (utilization exceeds 100 percent); $B = 6$ is the smallest stable batch and $B = 8$ the smallest that meets the 75 percent utilization target. Among the tabulated sizes, $B = 8$ has the lowest mean latency, and evaluating every integer batch size places the minimum of this model near $B = 10$. $B = 16$ has similar mean latency with lower utilization, while $B = 32$ leaves more headroom but pays a large accumulation-delay penalty.

Size B	$T_{\text{svc}}(B)$ (ms)	req/s	ρ_{serv}	$\mathbb{E}[W_q]$ (ms)	$(B-1)/(2\lambda_{\text{arr}})$ (ms)	$\mathbb{E}[t]$ (ms)
1	50.5	19.8	505% (unstable)	∞	0	∞
4	52	76.9	130% (unstable)	∞	15	∞
8	54	148.1	67.5%	56.1	35	145.1
16	58	275.9	36.3%	16.5	75	149.5
32	66	484.8	20.6%	8.6	155	229.6

Four operational regimes emerge across these batch sizes:

1. **Unstable regime (batch sizes 1 to 5)**: Utilization exceeds 100 percent, so the system cannot keep up with arrivals and queues grow without bound. The stability boundary of equation F.3, $B_{\text{min}} \approx 5.3$, falls inside this range.
2. **Utilization-target threshold (batch size 8)**: At 67.5 percent utilization, the system meets the 75 percent target, but queuing delay still contributes substantially to latency (56.1 ms). The steep occupancy trajectory charted in figure F.1 isolates why. Batch sizes between the stability boundary and this threshold push utilization past the $\rho_{\text{serv}} \approx 0.80$ hockey-stick knee into the caution band, where queuing delay surges nonlinearly as the curve approaches saturation.
3. **Latency knee (batch sizes 8 to 16)**: Mean latency is lowest in this range, with the model's minimum near $B = 10$. Batch size 16 trades a few milliseconds of additional total latency for lower utilization and queue wait time (16.5 ms vs. 56.1 ms), while $B = 32$ leaves still more utilization headroom but pays too much batch-accumulation delay.
4. **Diminishing returns (batch sizes beyond 32)**: Further batch size increases reduce utilization in this surrogate but increase accumulation delay. The separate KV-cache example below supplies a decode-memory bound; it is not part of this prefill-time calculation.

Applying Little's Law provides an internal consistency check for these results. Because $Q_{\text{req}} = \lambda_{\text{arr}} \cdot T_{\text{lat}}$ holds regardless of the arrival or service distribution, multiplying the arrival rate by the decomposed latency must recover the sum of requests in service and waiting in the queue.

Napkin Math F.1: Applying Little's Law to verify

Verification: By Little's Law (Little 1961), $Q_{\text{req}} = \lambda_{\text{arr}} \cdot T_{\text{lat}}$:
At $B = 16$ with $\lambda_{\text{arr}} = 100$ req/s and $\mathbb{E}[t] = 149.5$ ms:
$Q_{\text{req}} = 100 \times 0.1495$ s = 14.95 requests in system
With batch size 16 and utilization 36.3 percent:

- Expected requests in service: $16 \times 0.363 = 5.8$
- Expected requests in queue: 14.95 - 5.8 = 9.15

This matches the response-time decomposition: roughly 6 requests are in service and roughly 9 are waiting or accumulating into batches on average.

F.1.1.5 Batch size selection framework given an SLO

The 100 QPS analysis screened candidate batches by mean response time $\mathbb{E}[t]$, but production serving contracts enforce tail latency percentiles (for example, p99 latency < 200 ms) rather than averages. Selecting a production batch size under a tail-latency SLO follows a five-step procedure:

Step 1 characterize service time Measure T_{fixed} and t_{req} empirically by profiling inference at batch sizes 1, 8, and 32. Fit the linear model $T_{\text{svc}}(B) = T_{\text{fixed}} + t_{\text{req}} B$.

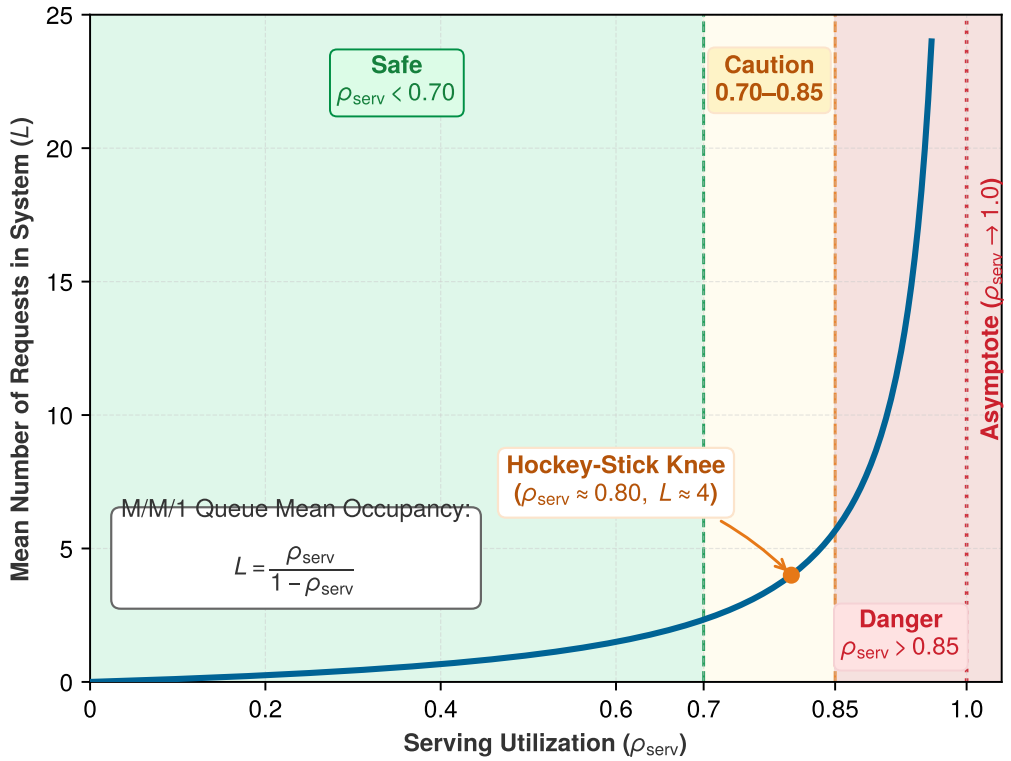

Figure F.1: The Queuing Hockey Stick: For an M/M/1 queue, $\rho_{serv}/(1-\rho_{serv})$ is the expected number of requests in the system, including the request in service. Its hyperbolic rise near saturation explains why production systems retain utilization headroom; the shaded operating bands are illustrative rather than universal thresholds.

Step 2 compute stability threshold Find the minimum batch size B_{min} such that $\mu_{eff}(B_{min}) > \lambda_{arr}$. Substituting the linear service-time model into $B/T_{svc}(B) = \lambda_{arr}$ yields the stability boundary:

$$B_{min} = \frac{T_{fixed}\lambda_{arr}}{1 - t_{req}\lambda_{arr}} \tag{F.3}$$

Choose an integer batch size strictly above this boundary, assuming $\lambda_{arr} t_{req} < 1$. Equality leaves no stability margin; if the denominator is nonpositive, batching alone cannot sustain the arrival rate.

Step 3 compute latency at candidate batch sizes Under the same Poisson-super-request surrogate and full-batch assumption, evaluate the expected response time across the candidate grid $B \in \{B_{min}, 2B_{min}, \ldots, B_{max}\}$. Summing the queuing, batch-accumulation, and execution terms gives the total expected latency:

$$\mathbb{E}[t(B)] = \frac{\lambda_{arr} \cdot T_{svc}(B)^2}{2B(1 - \lambda_{arr} T_{svc}(B)/B)} + \frac{B-1}{2\lambda_{arr}} + T_{svc}(B)$$

Step 4 account for tail latency If measured response times are approximately Gaussian, their p99 can be estimated from the mean and standard deviation. Adding 2.33 standard deviations captures the 99th percentile of a normal distribution:

$$T_{lat,p99} \approx \mathbb{E}[t] + 2.33 \cdot \sigma_t$$

where σ_t is the standard deviation of response time. This Gaussian rule is not a general heavy-traffic approximation because queuing response times are commonly skewed. For exponential response times, p99 is $-\ln(0.01)\mathbb{E}[t] \approx 4.6 \cdot \mathbb{E}[t]$; use measured percentiles for production sizing.

Step 5 select optimal batch size When the objective is maximum throughput under the increasing service-rate model, choose the largest stable, memory-feasible integer batch size whose p99 latency meets the SLO. Writing that constraint as an optimization problem over the candidate grid gives:

$$B = \max\{B : T_{\text{lat,p99}}(B) \leq \text{SLO}\}$$

Operating rules for batch size selection (table F.2) balance stability, latency SLOs, and cost efficiency. These rules generate candidate configurations, and load tests must then verify their latency, capacity, and cost under the actual request distribution:

Table F.2: Batch Size Decision Framework: Systematic approach to selecting batch size based on stability, latency SLO, and utilization targets.

Condition	Recommended action	Rationale
$B < B_{\text{min}}$	Increase batch size or add replicas	System is unstable
$T_{\text{lat,p99}} > \text{SLO}$	Re-evaluate stable batch sizes or add replicas	Latency exceeds target
$\rho_{\text{serv}} < 0.5$	Consider reducing replicas to save cost	System is overprovisioned
$\rho_{\text{serv}} > 0.85$	Add replicas for headroom	Approaching instability

F.1.1.6 Trade-off curves for visualizing the operating region

Batch execution sets two pressures against each other. As batch size B grows from 1 to 128, throughput climbs with diminishing returns as fixed overheads are amortized. Batch service time expands steadily, so a latency budget caps how much of the throughput gain a server can take (figure F.2).

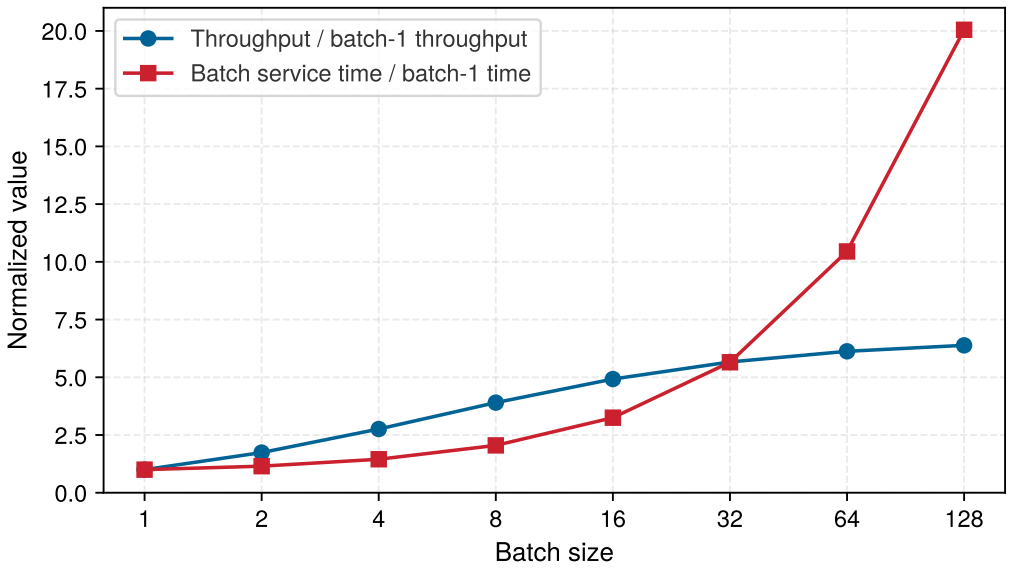

Figure F.2: Batch Size Trade-Off Curves: Dimensionless schematic with batch service time t(B) = 1 + 0.15(B – 1) and throughput B/t(B), each normalized to its batch-1 value. Larger batches amortize fixed work but take longer to complete. These hypothetical curves are not a benchmark, exclude queuing and batch-formation delay, and establish neither a memory-capacity boundary nor an optimal batch size.

Four properties of these curves guide the choice of batch size across candidate operating points. Together they explain why throughput and latency cannot be optimized independently:

1. **Pareto Frontier**: An operating point is Pareto efficient when no other candidate provides at least as much throughput and no more latency, with a strict improvement in at least one objective. The frontier is the set of these nondominated points.

2. **Knee of the Curve**: A knee marks diminishing throughput gains relative to increasing service time; it does not by itself select the best operating point.
3. **SLO-Constrained Optimum**: The best feasible batch size depends on the optimization objective and measured end-to-end latency, including queuing and batch formation. An SLO boundary can constrain that choice without guaranteeing an intersection or a unique optimum.
4. **Diminishing Returns**: In the schematic, fixed work becomes a smaller fraction of each batch, so successive batch-size increases yield smaller throughput gains.

The queuing-theoretic framework provides the mathematical foundation for the batching strategies introduced in section 10.6. Where intuition might suggest "larger batches are better for throughput," stability constraints, latency targets, and diminishing returns create a well-defined optimal operating region that varies by model type and deployment requirements. For autoregressive transformers, any candidate batch inside that queuing-feasible region must also fit its live attention state into accelerator memory.

F.1.2 KV cache fundamentals

During a single 2,048-token GPT-3 generation request, the key and value vectors accumulated across all 96 layers consume roughly 9.4 GB of high-bandwidth memory. A batch of 32 concurrent requests therefore exhausts the remaining memory of an eight-GPU A100 node before compute ever saturates. Section 10.7 defines the KV cache and the paged allocation that manages it, and section A.7 gives its per-token size. Sizing a serving node around those per-token states requires accounting for three coupled effects: how caching trades quadratic recomputation for linear memory growth, how the resulting footprint scales with sequence length and batch size, and why naive contiguous reservation fragments HBM (Kwon et al. 2023).

For a standard attention implementation at fixed model width, autoregressive generation without caching requires $\mathcal{O}(S^2)$ computation per token because each transformer layer recomputes the full prefix, including attention among its tokens. The KV cache stores previously computed key and value vectors, allowing each new token to attend to the prefix in $\mathcal{O}(S)$ computation. For serving at scale, this compute saving creates a memory-management challenge, because KV cache memory can exceed model weights for long contexts. The cache size grows linearly with both sequence length and batch size as calculated by equation F.4:

$$M_{\text{KV}} = 2 \times N_L \times d \times S \times B \times s_{\text{elem}} \tag{F.4}$$

For standard multihead attention, where the aggregate KV width equals the hidden width, N_L is the number of layers, d is the hidden dimension, S is sequence length, B is batch size, s_{elem} is storage size per element in bytes. The factor of 2 accounts for both keys and values. Grouped-query or multiquery attention instead uses the aggregate KV-head width in place of d.

Evaluating equation F.4 for a concrete multi-GPU node shows how quickly the KV cache bounds concurrency. For the 175B GPT-3 model ($N_L = 96$, $d = 12{,}288$, FP16 $s_{\text{elem}} = 2$ bytes) at its training context length $S = 2{,}048$, each request requires:

$$M_{\text{KV/req}} = 2 \times 96 \times 12288 \times 2048 \times 2 \text{ B} \approx 9.66 \text{ GB/request}$$

Across 8× A100 80-GB GPUs (687.2 GB total HBM), storing model weights in FP16 consumes 350 GB, leaving 337.2 GB for activation and KV cache memory. Dividing that residual capacity by the per-request cache gives a capacity-only upper bound of 34 concurrent requests. This bound assumes balanced sharding and reserves no memory for activations, workspaces, runtime state, or allocator losses; an implementable batch must be smaller (device nameplate capacities are converted to decimal GB before subtracting the decimal model footprint). Together with table F.1, this capacity ceiling closes the two-layer serving model: queuing dynamics set the minimum batch size needed for stability and latency headroom, while HBM capacity caps the maximum concurrent batch at long context lengths.

Summary

Serving performance is bounded by queues and by memory. As utilization ρ_{serv} approaches 1, queue depth and latency grow without bound, so production systems retain headroom suited

to their workload and SLO. Larger batches raise throughput but add batch-accumulation delay, which produces a latency knee rather than monotonic improvement; the best batch size depends on whether the objective is minimum mean latency, maximum SLO-compliant throughput, or cost. Little's Law, $Q_{\text{req}} = \lambda_{\text{arr}} T_{\text{lat}}$, gives a one-line consistency check on any serving measurement. The KV cache grows with sequence length, batch size, and KV-head width and can exceed the model weights (principle 18); paging reduces fragmentation and reservation losses without removing the underlying footprint.

G

System Assumptions

Every quantitative example in this book, from cluster failure rates to collective-communication costs to carbon footprint estimates, uses a specific set of assumed numbers: canonical cluster sizes, fabric bandwidths, accelerator reliability, power usage effectiveness, regional carbon intensity, and model utilization ranges. This appendix collects those fleet-scale assumptions in one place so the book's arithmetic can be audited, checked against chapter napkin math, or substituted into worked examples.

How to Use This Appendix

Use this appendix to verify a fleet-scale estimate, swap in an alternative assumption (for example, NDR vs. HDR fabric, Quebec vs. Poland carbon intensity, or a different model FLOPs utilization (MFU)), or identify which numerical constants a chapter calculation depends on. Locate the relevant section, read the `Value` and `Unit` columns, and substitute those values into target arithmetic:

- **When auditing or substituting the baseline 175B training job**: Use section G.1 and the Foundational Hardware Recap section.
- **When scaling a calculation across cluster sizes or network fabrics**: Use the Cluster Reference Configurations, Network and Interconnect Specifications, and Communication Model Parameters sections.
- **When estimating failure frequency or recovery downtime**: Use the component lifetime and recovery tables in the Reliability Assumptions section.
- **When pricing power, water, carbon, or cloud infrastructure**: Use the Sustainability Assumptions and Economic Assumptions sections.
- **When discounting peak FLOP/s to delivered throughput or checking units and sources**: Use the Capacity Planning Assumptions and Unit Conventions sections, and trace each constant's evidence class in section G.2.

Three worked estimates in napkin math G.1 combine these constants to calculate cluster failure rate, AllReduce time, and monthly carbon footprint. Hardware and link bandwidths are ceilings unless a chapter states an effective-bandwidth discount; treat the tables as this edition's shared assumption sheet rather than immutable physical laws.

Napkin Math G.1: Quick calculations with these assumptions

These tables support quick distributed-systems napkin math. Three examples show how several assumptions combine.

Problem: How many GPU failures should a large training cluster expect per day?

Variables: Use a cluster of 8,192 GPUs and an individual GPU mean time to failure (MTTF) of 50,000 hours.

Math: The GPU failure rate is 8,192 GPUs/50,000 hours $\approx$ 0.16 failures/hour, or roughly 3.9 failures/day.

Result: This estimate excludes network interface card (NIC), PSU, and cable failures. At 100,000 GPUs, the GPU-only rate is approximately 48 failures/day; adding NIC, PSU, and cable failures, whose MTTFs appear in table G.6, raises the operational incident stream further.

Systems insight: At fleet scale, failure becomes a continuous background process rather than an exceptional event.

Problem: How long does one AllReduce of the gradients of a 70B model take?

Variables: Use BF16 parameters and InfiniBand NDR with a nominal byte rate of 50 GB/s per port.

Math: The gradient payload is $70 \times 10^9 \times 2$ bytes = 140 GB. The ring AllReduce napkin estimate is 2×140 GB/50 GB/s ≈ 5.6 s.

Result: The collective can consume a significant fraction of a training step.

Systems insight: Pipeline and tensor parallelism exist partly to shrink the gradient volume each rank must communicate.

Problem: What is the monthly carbon footprint of a 10,000-GPU training run, and how much does the grid change it?

Variables: Each H100 draws 700 W at thermal design power (TDP). Use power usage effectiveness (PUE) 1.12, Quebec grid intensity 20 g/kWh, and Poland grid intensity 820 g/kWh.

Math: A GPU-only facility-power proxy is 10,000 GPUs $\times$ 700 W $\times$ 1.12 = 7.84 MW. In Quebec, 7.84 MW $\times$ 720 h/month $\times$ 20 g/kWh = 112.9 t CO_2 per month.

Result: The same run in Poland emits 4,628.7 t, a 41× difference.

Systems insight: Grid carbon intensity can dominate the emissions difference for the same hardware, model, and training duration. This proxy excludes host, memory, network, and storage power; total IT power must include those loads before applying PUE for a complete facility estimate.

Foundational Hardware Recap

Before a fleet calculation can attribute a slowdown to collective communication or stragglers, it needs the compute, memory-bandwidth, capacity, and thermal ceilings of one accelerator in isolation. The book's primary single-node reference is the NVIDIA H100 (table G.1), and Appendix A derives the single-machine quantities built on these specifications, including the ridge point, the memory hierarchy, and the per-parameter memory of training.

Table G.1: Single-Node H100 Assumptions: Peak specs from (Choquette 2023) and NVIDIA H100 documentation. These values provide the R_{peak} and BW baselines for iron-law calculations.

Category	Assumption	Value
Compute	H100 FP16 peak throughput	989 TFLOP/s
Compute	H100 FP8 peak throughput	1,979 TFLOP/s
Memory	HBM3 bandwidth	3.35 TB/s
Memory	HBM3 capacity	80 GB
Thermal	TDP	700 W

G.1 Reference Workload

SINGLE-ACCELERATOR ceilings turn into concrete multi-node requirements as soon as a model's training state exceeds the capacity of one H100. Where a chapter needs a concrete training job, it uses the Archetype A running instance that section 1.6.1.1 fixes (table G.2), so that any calculation built on that baseline can be audited or rerun with a different model or cluster. A chapter that studies a different model or cluster size says so explicitly.

Table G.2: Archetype A Running Instance: Mixed-precision training with the Adam optimizer (BF16 weights and gradients, FP32 master weights, two FP32 Adam moments); the checkpoint omits gradients. Section A.6 derives the per-parameter bytes. The reference cluster is the book's 1,024-GPU training cluster; the canonical cluster tiers below illustrate how behavior changes as scale departs from it.

Assumption	Value
Model parameters (dense transformer)	175B
Training state per parameter	16 bytes
Total training state	2.8 TB
Resumable checkpoint per parameter	14 bytes
Resumable checkpoint size	2.45 TB
Reference cluster	1,024 H100 GPUs (128 nodes of 8)

Cluster Reference Configurations

A communication or fault-tolerance design that works cleanly on the 1,024-GPU Archetype A reference cluster can be unnecessary on a 256-GPU research allocation yet fail completely on a 100,000-GPU hyperscale fleet. Four **canonical cluster tiers** (256, 2,048, 8,192, and 100,000 GPUs) therefore anchor how system behavior changes across scale (table G.3). They are editorial reference points aligned with common research-lab, production, large-training, and hyperscale fleets, not measurements from a single deployment, and chapters derive node counts, failure rates, and network requirements from these baselines and the assumptions in subsequent tables.

Table G.3: Canonical Cluster Sizes: Book-convention tier sizes for fleet-scale examples. At 8 per node: 32 nodes, 256 nodes, 1,024 nodes, and 12,500 nodes. Failure handling shifts from exception to steady state between the 256- and 8,192-GPU tiers (Kokolis et al. 2025).

Assumption	Value	Unit
Small cluster GPU count (256 tier)	256	GPUs
Medium cluster GPU count (2,048 tier)	2,048	GPUs
Large cluster GPU count (8,192 tier)	8,192	GPUs
Mega cluster GPU count (100,000 tier)	100,000	GPUs

The relationship between cluster size and system behavior is not linear. A 256-GPU cluster might go days between failures; a 100,000-GPU cluster experiences multiple failures per hour. This nonlinearity is why chapter 7 and chapter 8 treat failure as a continuous background process rather than an exceptional event. The network fabric determines whether gradient synchronization across those nodes takes milliseconds or seconds and whether pipeline parallelism is viable across node boundaries.

Network and Interconnect Specifications

Distributed training and inference are bottlenecked by data movement between machines as often as by computation within them. The network assumptions fix bandwidth and latency for every interconnect tier, from NVLink/PCIe within a node (NVIDIA Corporation 2017, 2020; Choquette 2023), through InfiniBand, remote direct memory access (RDMA), and Ethernet fabrics (InfiniBand Trade Association 2000; Gangidi et al. 2024; NVIDIA 2026f; Ethernet Alliance 2025), to the speed-of-light WAN floor (physics identity).

Intra-node interconnects

Splitting a single transformer layer across eight GPUs inside a server requires an all-reduce on every forward and backward pass, which is viable at hundreds of gigabytes per second over NVLink but stalls over a host PCIe bus. Those intra-node bandwidths and latencies (table G.4) determine which fine-grained parallelism strategies must remain within the node boundary.

Table G.4: Intra-Node Interconnect Bandwidth and Latency: NVLink/PCIe peaks from vendor documentation (NVIDIA Corporation 2017, 2020; Choquette 2023). The displayed NVLink rates are bidirectional aggregates; PCIe rates are per direction. Comparing like-for-like per-direction rates gives an NVLink/PCIe ratio of 7–10×, informing placement of communication-intensive tensor parallelism.

Assumption	Value	Unit
NVLink bandwidth (V100)	300	GB/s
NVLink bandwidth (A100)	600	GB/s
NVLink bandwidth (H100)	900	GB/s
PCIe Gen3 bandwidth (V100 host)	15.75	GB/s
PCIe Gen4 bandwidth (A100 host)	32	GB/s
PCIe Gen5 bandwidth (H100 host)	64	GB/s
NVLink one-way latency	500	ns
PCIe Gen5 one-way latency	1000	ns

Inter-node interconnects

Once data crosses the node boundary, bandwidth drops by an order of magnitude and latency increases by 5–10×. This cliff shapes every decision about parallelism strategy. Operations that require frequent, fine-grained communication (tensor parallelism) stay within the node, while operations with coarser communication patterns (data parallelism, pipeline parallelism) span nodes. To accommodate the different conventions used across chapters, table G.5 catalogs both the bit-rate and byte-rate forms of each fabric.

Table G.5: Inter-Node Network Bandwidth and Latency: Fabric link rates from (InfiniBand Trade Association 2000) and product-line specifications. The 9× NVLink-vs-NDR gap explains topology-aware parallelism (chapter 5, chapter 6).

Assumption	Value	Unit
InfiniBand HDR link bandwidth (bit rate)	200	Gb/s
InfiniBand HDR bandwidth (byte rate)	25	GB/s
InfiniBand NDR link bandwidth (bit rate)	400	Gb/s
InfiniBand NDR bandwidth (byte rate)	50	GB/s
InfiniBand XDR bandwidth (byte rate)	100	GB/s
400 GbE bandwidth (byte rate)	50	GB/s
800 GbE bandwidth (byte rate)	100	GB/s
100 GbE RoCE bandwidth (byte rate)	12.5	GB/s
InfiniBand one-way latency	5000	ns

Wide-area network

Cross-region communication introduces latency floors set by the speed of light in optical fiber. For a 5,000 km link (roughly New York to London), the minimum one-way latency is 5,000 km/200,000 km/s = 25 ms, orders of magnitude higher than intra-data-center latency. This physical constraint is why geo-distributed training uses asynchronous methods or careful placement of pipeline stages, and why federated learning (chapter 11) tolerates stale gradients. Network specifications define how fast data moves across links and regions, but a multi-thousand-node training job can sustain those transfer rates only while every participating accelerator, network card, and switch remains healthy.

Reliability Assumptions

At fleet scale, component failures are statistical certainties rather than exceptional events. The reliability assumptions here fix MTTF for each major component class in a data-center GPU node and the time to detect and recover from failures. These values drive analysis in chapter 7 and scheduling in chapter 8. Deeper failure modeling appears in Appendix E.

Component MTTF

A single data-center GPU runs for tens of thousands of hours on average before failing, yet combining eight GPUs and eight NICs per server across a thousand nodes turns multi-year component lifetimes into a cluster interruption every few hours. The steady-state component MTTFs in table G.6 exclude the early "infant mortality" period after deployment and the late "wear-out" period beyond rated lifetime, drawing on large-GPU research-cluster analysis (Kokolis et al. 2025), TPUv4 supercomputer resiliency and operations (Zu et al. 2024), and warehouse-scale machine design (Barroso et al. 2019).

Table G.6: Component Mean Time to Failure: Steady-state MTTF anchors from (Kokolis et al. 2025; Zu et al. 2024; Barroso et al. 2019). Combined DGX H100 node MTBF ≈ 4,687.5 h (195.3 d) for 8 GPUs and 8 NICs in series, the same node composition as the MTBF cascade in section E.1.3. The node's 6 PSUs run 4+2 redundant, so no single PSU failure stops it and they add no term. The PCIe switch and HBM rows are node components that series model leaves out, and the top-of-rack switch and cable rows are network failures outside the node. Counting any of them would shorten the interval between interruptions a job sees.

Assumption	Value	Unit
GPU MTTF	50,000	hours
NIC MTTF	150,000	hours
PSU MTTF	100,000	hours
PCIe switch MTTF	200,000	hours
HBM MTTF	200,000	hours
Top-of-rack switch MTTF	300,000	hours
Optical cable and transceiver MTTF	50,000	hours

Recovery time assumptions

When one of those components fails, the cluster loses wall-clock time not only to replaying steps since the last snapshot, but also to expiring the heartbeat timer, allocating a replacement node, and streaming the checkpoint back from storage. Table G.7 lists the detection, rescheduling, and storage-bandwidth assumptions used in checkpoint interval optimization (the Young-Daly formula in section 7.8.1) and effective throughput calculations throughout this book.

Table G.7: Recovery Time Parameters: Design assumptions for heartbeat, reschedule, and checkpoint-write bandwidth (Young-Daly context (Young 1974; Daly 2006)). Recovery before replay ≈ detect + reschedule + checkpoint size divided by storage bandwidth; for a 70B-parameter model with a 14-byte-per-parameter checkpoint, about 99.8 s per failure.

Assumption	Value	Unit
Failure detection timeout (heartbeat)	30	seconds

Assumption	Value	Unit
Job reschedule delay	60	seconds
Checkpoint write bandwidth (aggregate)	100	GB/s

Together with table G.6, these recovery parameters determine how much wall time is lost when hardware fails. Even on fault-free steps, however, distributed training stalls whenever workers synchronize gradients or activations across the network.

Communication Model Parameters

A small control message or fine-grained collective pays a per-message software and fabric launch delay on every hop even on a 400 Gbps link, whereas a multi-gigabyte gradient synchronization is governed by sustained byte rate. The book's communication models capture both regimes with the α-β framework, $T(n) = \alpha + n/\beta$, where α is startup latency and β is sustained bandwidth (Hockney 1994). Table G.8 lists values from fabric specifications, vendor product documentation, and large-cluster training studies (InfiniBand Trade Association 2000; Gangidi et al. 2024; NVIDIA 2026f; Ethernet Alliance 2025; Z. Jiang et al. 2024), feeding the collective-communication models in chapter 5 and parallelism comparisons in chapter 6.

Table G.8: Communication Model Parameters (α-β): α (startup) and β (bandwidth) from fabric specifications and measured cluster studies (InfiniBand Trade Association 2000; Z. Jiang et al. 2024). Every α is an effective per-message launch latency that includes the communication library's software stack, the convention of table 5.4, so it differs from the fabric's end-to-end one-way latency in table G.5. The InfiniBand NDR value is the single anchor used in the book's collective models. The 33× NDR-vs-TCP gap explains why RDMA fabrics dominate synchronous training.

Assumption	Value	Unit
InfiniBand NDR startup latency (α)	1.5	us
InfiniBand NDR sustained bandwidth (β)	50	GB/s
InfiniBand HDR startup latency (α)	2–5	us
InfiniBand HDR sustained bandwidth (β)	25	GB/s
RoCE startup latency (α)	5–10	us
RoCE sustained bandwidth (β)	12.5	GB/s
TCP startup latency (α)	50	us

AllReduce cost model

Synchronizing a M-byte gradient buffer across N data-parallel ranks via ring AllReduce transmits $2(N-1)M/N$ bytes per rank, which approaches $2M$ bytes per rank as N grows. Table G.9 lists the asymptotic bandwidth factor and node-boundary count used in AllReduce cost estimation.

Table G.9: AllReduce Cost Parameters: Ring AllReduce factor of 2 from the standard ring algorithm (Gibiansky 2017). The DGX H100 node (8 GPUs per node) sets the intra-node versus inter-node boundary in hierarchical collectives.

Assumption	Value	Unit
Ring AllReduce bandwidth factor	2	(dimensionless)
GPUs per node	8	GPUs/node

Compute, network, and reliability constants determine how long a training or serving job runs on the cluster; operating that cluster in a real data center also requires budgeting the electrical power, cooling water, and grid carbon consumed by the surrounding facility.

Sustainability Assumptions

The environmental impact of fleet-scale ML depends on three factors: how much power the facility draws beyond the IT equipment (PUE), how much water the cooling system consumes (water usage effectiveness (WUE)), and how much carbon the local grid emits per kilowatt-hour. The sustainability assumptions here support chapter 15 and total-cost-of-ownership estimates in chapter 12.

Power usage effectiveness

A 10,000-GPU cluster drawing megawatts of IT power requires additional megawatts for the chillers, fans, and power distribution that keep the accelerators within their thermal envelope. That facility overhead is captured by PUE, the ratio of total facility power to IT equipment power (equation 15.2), where a PUE of 1.0 would mean zero energy beyond the IT equipment itself. The reference tiers in table G.10 represent illustrative facility scenarios rather than universal bounds of a cooling technology.

Table G.10: Power usage effectiveness tiers: Illustrative PUE tiers from hyperscale surveys and Green Grid definitions (Uptime Institute 2022; The Green Grid 2007). Industry-average (2023) vs. liquid-cooled gap: 49.1 percent more facility power for the same IT load.

Assumption	Value	Unit
PUE (liquid-cooled, best in class)	1.06	(ratio)
PUE (best air-cooled)	1.12	(ratio)
PUE (industry average, 2022)	1.55	(ratio)
PUE (industry average, 2023)	1.58	(ratio)

Water usage effectiveness

Evaporative cooling towers achieve low PUE by evaporating water to reject heat, whereas a closed liquid loop limits water loss at the rack while still depending on the heat-rejection system outside that loop. To quantify that cooling trade-off, WUE measures liters of water consumed per kilowatt-hour of IT energy (The Green Grid 2011), and table G.11 separates air-cooled, evaporative, and closed-loop liquid reference tiers.

Table G.11: Water usage effectiveness tiers: Illustrative WUE tiers for cooling-class comparisons (The Green Grid 2011; Uptime Institute 2022). Evaporative example at 10 MW IT load: 432,000 L/day.

Assumption	Value	Unit
WUE (air-cooled)	1.8	L/kWh
WUE (evaporative cooling)	1.8	L/kWh
WUE (closed-loop liquid)	0	L/kWh

Regional carbon intensity

The same training run emits vastly different amounts of CO_2 depending on the grid that supplies its electricity. Table G.12 lists regional grid carbon intensities from IEA and Canadian government emissions-factor data (International Energy Agency 2024b; Environment and Climate Change Canada 2026), rounded for napkin math. These values are central to location-aware scheduling in chapter 15.

Table G.12: Regional Grid Carbon Intensity: Grid carbon intensity in gCO_2/kWh, anchored to emissions-factor data (International Energy Agency 2024b; Environment and Climate Change Canada 2026) and ordered from low (Norway) to high (Poland). Location alone changes carbon intensity by 82×, often more than an algorithmic change saves.

Assumption	Value	Unit
Grid carbon intensity (Norway)	10	gCO2/kWh
Grid carbon intensity (Quebec)	20	gCO2/kWh
Grid carbon intensity (France)	50	gCO2/kWh
Grid carbon intensity (EU average)	270	gCO2/kWh
Grid carbon intensity (US average)	429	gCO2/kWh
Grid carbon intensity (Poland)	820	gCO2/kWh

Power density

Packing eight high-wattage accelerators and their host and network hardware into a server chassis pushes AI racks far past the thermal load of conventional enterprise racks. Whether liquid cooling is required depends on the installed heat-removal capacity, airflow, and server design rather than a universal rack-power threshold, and table G.13 lists the reference rack power levels used in chapter 2 and chapter 15, anchored to current AI power-and-cooling product disclosures (Dell Technologies 2026; Schneider Electric 2024).

Table G.13: Rack Power Density: Illustrative rack-power tiers and air-cooling limit (Uptime Institute 2022; Dell Technologies 2026; Schneider Electric 2024). The listed AI rack scenarios (70 kW–100 kW) exceed the illustrative 30 kW air-cooling reference; the feasible limit depends on the installed system.

Assumption	Value	Unit
Rack power (traditional enterprise)	12	kW
Rack power (AI, typical)	70	kW
Rack power (AI, high density)	100	kW
Air-cooling practical limit per rack	30	kW

These sustainability constants quantify the megawatts, water volume, and carbon emissions of operating a fleet; converting those physical resource draws and accelerator allocations into project budgets requires unit prices for compute hours, electricity, and network egress.

Economic Assumptions

A multi-week training run on thousands of GPUs or a global serving tier streaming responses across regions accumulates millions of dollars in compute rental, facility electricity, and network egress charges. The prices in table G.14 provide illustrative hyperscaler-order rates (2024–2025 magnitude) for total cost of ownership (TCO) ratio analysis in chapter 12 and chapter 10, not quotes for a specific contract.

Table G.14: Fleet-Scale Economic Parameters: Illustrative GPU-hour, electricity, and egress rates for TCO examples (substitute contract pricing for absolute budgets). The fleet GPU-hour reference is the price per in-service accelerator-hour used by this book's fleet cost examples; the cloud rental rates are public list-price anchors for comparison.

Assumption	Value	Unit
Fleet GPU-hour reference price	2	dollar/h
Cloud GPU training rental rate	4	dollar/h
Cloud GPU inference rental rate	2.5	dollar/h
Cloud electricity price	0.12	dollar/kWh
Cloud egress price per GB	0.09	dollar/GB

Pricing an accelerator-hour settles what a cluster allocation costs per unit of wall time, but converting that hourly spend into delivered training progress requires discounting datasheet peak FLOP/s by kernel utilization, multi-node scaling efficiency, and operational downtime.

Capacity Planning Assumptions

Multiplying a cluster's GPU count by the H100's datasheet peak FLOP/s overstates delivered training throughput by more than a factor of three, because single-device kernel inefficiency, multi-node communication stalls, and out-of-step operational downtime each take a cut. The capacity-planning assumptions quantify the fraction of peak hardware performance that real workloads sustain on one device (MFU), how efficiently that performance scales as more machines are added (scaling efficiency), and what share of wall time is lost to operational pauses. Because these factors compound across the fleet, the planning equation is multiplicative:

$$R_{\text{eff}} = N R_{\text{peak}} \times \text{MFU} \times \eta_{\text{scaling}} \times (1 - f_{\text{overhead}})$$

Here R_{eff} is effective fleet throughput, $N R_{\text{peak}}$ is nominal cluster peak throughput. MFU is the single-device baseline measured over productive steps. η_{scaling} captures the additional in-step losses from distribution, including communication and pipeline bubbles. f_{overhead} removes only out-of-step checkpoint, recovery, and maintenance time. These windows must be disjoint: measured fleet-step MFU already includes scaling losses, so it must not be multiplied by scaling efficiency again. The three subsections that follow fix those factors in order.

Model FLOPs utilization

Even on a single accelerator, memory-bandwidth stalls, kernel launch gaps, and non-matmul operations prevent a model from sustaining datasheet peak throughput. MFU measures useful model FLOP/s divided by hardware peak FLOP/s over a stated measurement window, so an MFU of 0.50 means the workload achieves half of peak FLOP/s; synchronization and idle time within that window lower MFU even though they contribute no model FLOPs. Table G.15 lists the reference MFU ranges used in training time and cost estimates throughout this book.

Table G.15: Model FLOPs utilization ranges: Training MFU 0.3–0.5 from large-model training reports (Chowdhery et al. 2022; Narayanan et al. 2021); inference batch-1 MFU 5 percent from (Pope et al. 2023) (memory-bound decode).

Assumption	Value	Unit
Training MFU (low end)	0.3	(fraction)
Training MFU (high end)	0.5	(fraction)
Inference MFU (batch size 1)	0.05	(fraction)

Assumption	Value	Unit
Inference MFU (batched)	0.4	(fraction)

Scaling efficiency

Expanding a job from one node to hundreds of nodes adds gradient synchronization and pipeline bubble stalls that stretch each step beyond its single-device baseline. Scaling efficiency $\eta_{\text{scaling}} = T_1/(N \times T_N)$ measures how much of the added compute actually reduces training time, and table G.16 lists the tiers this book plans with. From 256 GPUs up, these values match MLSysIM's modeled efficiency curve for the 175B reference job at a fixed GPT-3 global batch. They extend the multi-GPU scaling sweep of figure 8.12 in chapter 8 out to 8,192 accelerators. The 32-GPU tier is a book convention, since one tensor- and pipeline-parallel replica of that job needs 64 GPUs. Production systems retain more at scale by training with a larger global batch and, where they pipeline, by interleaving stages. Published large language model (LLM) training runs built this way report near-linear scaling on thousands of accelerators (Z. Jiang et al. 2024; Chowdhery et al. 2022).

Table G.16: Scaling Efficiency by Cluster Size: Planning η_{scaling} tiers. The 256- to 8,192-GPU values are MLSysIM's model of the 175B reference job at a fixed global batch, and the 32-GPU value is a book convention. At 8,192 GPUs, 4,587.5 effective GPUs of 8,192 nominal.

Assumption	Value	Unit
Scaling efficiency (32 GPUs)	0.9	(fraction)
Scaling efficiency (256 GPUs)	0.87	(fraction)
Scaling efficiency (1,024 GPUs)	0.83	(fraction)
Scaling efficiency (8,192 GPUs)	0.56	(fraction)

Overhead budgets

Real training jobs spend a fraction of wall time on noncompute activities: pipeline bubble idle time, checkpoint writes, failure recovery, and scheduled maintenance. Table G.17 lists the overhead budgets assumed throughout this book. When measured as nonoverlapping fractions of the same wall-time window, these fractions are additive: total overhead at fleet scale is approximately 5 percent + 3 percent + 10 percent + 5 percent = 23 percent, leaving 77 percent outside these listed overhead categories. This combined budget includes pipeline bubbles; the multiplicative planning equation excludes them from its operational factor f_{overhead} because scaling efficiency η_{scaling} already accounts for them.

Table G.17: Overhead Budgets (Fraction of Wall Time): Book-convention wall-time fractions for pipeline bubbles, checkpointing, failure recovery, and maintenance, engineering targets, not physical constants. Combined 23 percent overhead ⇒ 23.1 days of effective work in a 30-day run.

Assumption	Value	Unit
Wall-time overhead (pipeline bubbles)	0.05	(fraction)
Wall-time overhead (checkpointing)	0.03	(fraction)
Wall-time overhead (failure recovery)	0.1	(fraction)
Wall-time overhead (scheduled maintenance)	0.05	(fraction)

Combining hardware peaks, link rates, reliability intervals, facility power, dollar rates, and utilization fractions in a single calculation requires a consistent set of dimensional prefixes and unit aliases.

Unit Conventions

Combining an accelerator memory capacity in gigabytes, a checkpoint payload in terabytes, and a link rate in gigabits per second produces silent factor-of-eight or base-2-versus-base-10 errors unless every table shares one dimensional system. Every quantitative example in this book therefore follows the unit conventions in table G.18, where each row gives the multiplier k in `1 alias` $= k$ `base`. Data prefixes use decimal SI (KB $= 10^3$ bytes, not 1024). Throughput quantities (FLOP/s, GB/s) combine these aliases with time; adding incompatible dimensions (bytes to FLOP/s) is a category error in napkin math, not a unit conversion.

Hardware capacity labels are the one deliberate exception. Reference tables may retain vendor nameplate numbers and labels; the byte budgets used in calculations follow the device assumptions stored in the book's hardware registry. For example, the book models an 80 GB accelerator with 80×2^{30} bytes, so a node of eight such devices holds 640×2^{30} bytes rather than 640×10^9. These are modeled capacity assumptions, not universal guarantees inferred from the labels. Every other data quantity, including model footprints, KV-cache sizes, checkpoint volumes, and collective payloads, uses those decimal SI prefixes. Where a worked example subtracts a computed footprint from a device capacity, the stored capacity byte budget is converted to decimal units alongside the footprint so that the arithmetic printed on the page is the arithmetic a reader can reproduce.

Table G.18: Unit Conventions: Decimal SI aliases and scale factors (book convention). Throughput forms (FLOP/s, GB/s) divide work or data by time using the same conventions.

Alias	Multiplier	Base unit
`byte`	1	byte
`KB`	1000	byte
`MB`	1e+06	byte
`GB`	1e+09	byte
`TB`	1e+12	byte
`PB`	1e+15	byte
`flop`	1	flop
`GFLOPs`	1e+09	flop
`TFLOPs`	1e+12	flop
`ZFLOPs`	1e+21	flop
`param`	1	param
`Mparam`	1e+06	param
`Gbps`	1e+09	bit/s
`NS`	1e-09	second
`US`	1e-06	second
`MS`	0.001	second
`second`	1	second
`hour`	3600	second
`day`	86400	second
`joule`	1	joule
`watt`	1	watt
`meter`	1	meter

Two fleet-scale conventions deserve emphasis beyond the table because they recur in almost every cross-node or facility calculation:

- Network bandwidth appears in both bit-rate (Gbps) and byte-rate (GB/s). Divide by 8 to convert Gbps to GB/s (for example, 400 Gb/s = 50 GB/s for InfiniBand NDR). Chapters use whichever form fits the calculation; byte-rate is usually easier for transfer-time estimates.
- Carbon intensity uses gCO_2/kWh (grams of CO_2 per kilowatt-hour). Multiply total energy (kWh) by regional carbon intensity to get emissions in grams, then divide by 10^6 for metric tons.

G.2 Assumption Provenance

The reference tables in this appendix mix H100 datasheet ceilings, empirical failure rates from warehouse-scale fleets, modeled MLSysIM scaling curves, and illustrative cloud rental rates, and treating all four as interchangeable measurements would misstate what an estimate proves. Each section therefore carries an explicit source classification and primary bibliography in table G.19, separating empirical hardware measurements from modeled scaling profiles and editorial conventions so a reader can see which numbers change with a hardware generation and which reflect a deployment choice.

Table G.19: Fleet Assumption Provenance Catalog: Quick map from appendix section to source class and bibliography.

Appendix section	Source type	Primary references
H100 recap	Vendor datasheet peak	(Choquette 2023)
Reference workload (Archetype A)	Book convention	GPT-3-class dense model; byte widths from the notation conventions
Cluster tiers	Book convention	(Kokolis et al. 2025) (scale context)
Network fabrics	Vendor specs; InfiniBand standard	(NVIDIA Corporation 2017, 2020; Choquette 2023; InfiniBand Trade Association 2000)
Reliability (MTTF)	Fleet studies	(Kokolis et al. 2025; Zu et al. 2024; Barroso et al. 2019)
Recovery times	Design assumptions	(Young 1974; Daly 2006)
Communication (α-β)	Fabric specs; training at scale	(InfiniBand Trade Association 2000; Z. Jiang et al. 2024)
AllReduce model	Algorithm identity	(Gibiansky 2017)
Sustainability (PUE/WUE/rack power)	Industry surveys	(Uptime Institute 2022; The Green Grid 2007)
Carbon intensity	Grid statistics	(International Energy Agency 2024b; Environment and Climate Change Canada 2026)
Cloud economics	Illustrative rates	Editorial (2024–2025 order of magnitude)
MFU	Published LLM training	(Chowdhery et al. 2022; Narayanan et al. 2021; Pope et al. 2023)
Scaling efficiency	MLSysIM model; book convention	Context: (Z. Jiang et al. 2024; Chowdhery et al. 2022)
Overhead budgets	Book convention	Editorial engineering targets
Unit conventions	Decimal SI; book notation	Editorial

Summary

Fleet-scale back-of-envelope estimates stay comparable across chapters only when they draw from a single shared assumption sheet. Every napkin-math calculation in this book traces to a row in the reference tables of this appendix, and table G.19 records whether that row comes from a datasheet, a fleet study, an MLSysIM model, an illustrative rate, or a book convention. Reliability rows set the failure clock, with failure a steady-state condition at 10,000 or more GPUs (Appendix E models it in depth). Communication rows decide whether synchronous training is viable, and the 33× startup-latency gap between InfiniBand NDR and TCP explains why training clusters favor RDMA fabrics. Sustainability rows expose location as a lever, since grid carbon intensity differs by 82× between Norway and Poland. Capacity rows bridge peak hardware to useful throughput. At 8,192

GPUs, about 56 percent of ideal scaling is realized for the reference job at a fixed batch, and the listed overheads consume about 23 percent of wall time. When a deployment differs in fabric, region, MFU, or cluster size, substitute its values and re-evaluate the corresponding formulas.

G

Glossary

This glossary defines key terms used throughout this book. Terms are organized alphabetically.

1

1F1B scheduling One-Forward-One-Backward scheduling, a synchronous pipeline scheduling algorithm that alternates forward and backward passes for in-flight microbatches, capping peak activation memory to $\mathcal{O}(N_{\text{PP}})$ where N_{PP} is pipeline depth.

3

3D parallelism The composition of data parallelism, pipeline parallelism, and tensor parallelism across orthogonal axes of a cluster mesh to match physical bandwidth tiers (intra-node NVLink, inter-node fabric, cluster spine).

A

accountability The mechanisms by which individuals or organizations are held responsible for the outcomes of AI systems, involving traceability, documentation, auditing, and the ability to remedy harms.

activation checkpointing A memory-saving training technique that stores only selected activations and recomputes others during backpropagation, trading extra compute for lower memory use.

active optical cable (AOC) A fiber-optic cable assembly with integrated electro-optical transceivers at both ends that converts high-speed electrical signals to optical pulses, providing reaches up to 100 meters across server racks.

adapter modules Small trainable neural network components inserted between frozen layers of a pretrained model to enable lightweight adaptation without modifying the base architecture.

adaptive routing A routing mechanism that selects among available paths using current network conditions, such as queue occupancy, to avoid hot spots while preserving the transport's ordering requirements.

advanced packaging Semiconductor integration techniques, such as 2.5D interposers, chiplets, and 3D stacking, that connect multiple dies with shorter, denser links than a conventional circuit board can provide.

adversarial attack A type of attack where carefully crafted inputs are designed to cause machine learning models to make incorrect predictions while remaining nearly indistinguishable from legitimate data to humans.

adversarial example A maliciously modified input that is designed to fool a machine learning model into making an incorrect prediction, often created by adding small, imperceptible perturbations to legitimate data.

adversarial training A defense technique that involves training models on adversarial examples to improve their robustness and ability to correctly classify adversarial inputs.

algorithmic fairness The principle and mathematical formulation that algorithmic decisions should not disproportionately or unjustly disadvantage individuals or protected groups based on sensitive demographic attributes.

alignment tax The measured degradation in general task performance, training throughput, or inference latency incurred as a byproduct of applying safety and alignment interventions.

AllGather A collective communication operation that gathers each rank's shard and makes the combined result available to every rank.

AllReduce A collective communication operation that reduces values across ranks and returns the reduced result to every rank, commonly used for gradient synchronization.

AllToAll A collective communication operation where each rank sends distinct data to every other rank and receives distinct data from each rank.

alpha-beta model A communication cost model that separates fixed startup latency from message-size-dependent transfer time to estimate the cost of moving data.

Amdahl's law The bound on the speedup of fixed work across N processors when a fraction r of it is serial, $1/(r + (1 - r)/N)$, which can never exceed $1/r$. At fleet scale, the serial fraction includes exposed communication, synchronization, and pipeline bubbles.

anomaly detection The identification of patterns in data that do not conform to expected behavior, often used to detect outliers, faults, or malicious activities in systems.

anonymization The process of removing or modifying personally identifiable information from datasets to protect individual privacy, though often insufficient against sophisticated re-identification attacks.

Archetype A (GPT-4/Llama-3) The book's recurring large dense language-model workload, used to probe communication-heavy distributed training and memory-bandwidth-bound serving. Its running instance is a 175B-parameter dense transformer trained with mixed precision and the Adam optimizer on 1,024 H100 GPUs.

Archetype B (DLRM at Scale) The book's recurring recommendation workload, a deep learning recommendation model whose large embedding tables are sharded across many nodes. It probes sparse feature routing, AllToAll communication, and tail latency.

Archetype C (Federated MobileNet) The book's recurring edge workload, an efficient convolutional network trained by federated learning across phone-class devices with sustained power budgets of a few watts. It probes the per-device compute envelope and the coordination of unreliable participants.

arithmetic intensity The amount of computation performed per byte of data moved, usually measured in FLOP/byte.

artificial intelligence A broad field of computer science focused on creating systems that can perform tasks typically requiring human intelligence, including learning, reasoning, and decision-making.

asynchronous checkpointing A fault-tolerance mechanism that snapshots model state into local staging memory before asynchronously streaming bytes to persistent storage in the background, minimizing training stall time.

attained service scheduling A duration-agnostic priority queueing discipline where job priority decreases monotonically as cumulative resource consumption increases, approximating shortest remaining processing time without requiring runtime estimates.

autoencoder A neural network architecture that learns compressed data representations by minimizing reconstruction error, commonly used for anomaly detection and dimensionality reduction.

automation bias The tendency for humans to over-rely on automated system outputs even when clear errors are present, potentially compromising human oversight.

AutoML Automated machine learning that uses algorithms to automate the process of applying machine learning to real-world problems, including feature engineering, model selection, and hyperparameter tuning.

availability The property that a storage service responds when a client needs it; it is distinct from whether acknowledged data survives failures or whether a read returns the newest committed version.

availability attack A type of data poisoning attack that aims to degrade the overall performance of a machine learning model by introducing noise or corrupting training data across multiple classes.

B

backdoor attack A malicious machine learning attack where a model is trained or poisoned to perform normally on standard inputs but produce adversary-selected behaviors whenever a specific secret trigger pattern is present.

backfill scheduling An optimization policy that schedules smaller, shorter queued jobs out of order onto currently idle resources, provided they do not delay the planned start time of higher-priority gang-scheduled jobs.

backpressure A control response that slows admission or bounds in-flight work when a downstream stage cannot sustain the upstream arrival rate.

backpropagation The algorithm for computing gradients in neural networks by propagating error signals backward through layers, essential for training but computationally expensive on resource-constrained devices.

bandwidth hierarchy The ordered set of data-movement tiers within and between computing systems, from on-chip memory through device, node, rack, and pod interconnects, whose bandwidth generally decreases as physical distance increases.

batched gather-matrix-vector (BGMV) A GPU kernel operation that batches matrix-vector products across heterogeneous adapters, gathering adapter-specific weights according to each token's adapter index.

Bayesian neural networks Neural networks that incorporate probability distributions over their weights, enabling uncertainty quantification in predictions and more robust decision-making.

bias detection Systematic methods for identifying unfair discrimination or disparate treatment across different demographic groups in machine learning system outputs.

bias mitigation Techniques and interventions designed to reduce unfair discrimination in machine learning systems, applied during data collection, model training, or postprocessing stages.

bias-only adaptation A lightweight training strategy that freezes all model weights and updates only scalar bias terms, drastically reducing memory requirements and computational overhead for on-device learning.

bin packing The assignment of jobs with resource demands into fixed-capacity nodes or clusters; in ML fleets this often spans GPUs, memory, CPUs, network locality, and topology.

bisection bandwidth The minimum aggregate link capacity crossing any cut that splits a network into two equal halves, bounding worst-case throughput for global communication.

bit flip A hardware fault where a single bit in memory or a register unexpectedly changes its value from 0 to 1 or vice versa, potentially corrupting data or computations.

black-box attack An adversarial attack where the attacker has no knowledge of the model's internal architecture, parameters, or training data, and must rely solely on querying the model and observing outputs.

block-wise quantization A quantization method that partitions tensors into independent sub-blocks with individual scaling factors, confining extreme outlier values to single blocks.

built-in self-test (BIST) Hardware testing mechanisms that allow components to test themselves for faults using dedicated circuitry and predefined test patterns.

bulk synchronous parallel (BSP) A parallel execution model in which workers perform local computation, communicate, and then wait at a global barrier before the next step.

C

C³ taxonomy The book's diagnostic framework for distributed execution, which divides wall-clock time and bottlenecks into Compute (C_1, local arithmetic and memory on each node), Communication (C_2, data movement between devices), and Coordination (C_3, synchronization, scheduling, and recovery). The Fleet Law assigns one term to each dimension.

canary deployment A rollout pattern that sends a small slice of traffic to a new model or service version, monitors behavior, and then expands or rolls back.

CAP theorem A distributed-systems result stating that, under a network partition, a system must choose between strict consistency and availability.

carbon-aware scheduling A workload scheduling strategy that shifts compute tasks geographically or temporally to match periods of high renewable energy availability and low grid carbon intensity.

carbon footprint The total amount of greenhouse gas emissions produced directly and indirectly by an individual, organization, event, or product, typically measured in CO_2-equivalent.

carbon intensity The greenhouse gas emissions associated with a unit of electricity, commonly measured in grams of CO_2-equivalent per kilowatt-hour.

carbon usage effectiveness (CUE) A data-center sustainability metric defined as the ratio of total carbon dioxide equivalent emissions produced by data center operations to total IT equipment energy consumption.

catastrophic forgetting The tendency of a neural network to abruptly lose proficiency on previously learned tasks or historical distributions when sequentially trained on new data without rehearsal or regularization.

certified robustness A formal guarantee that a model prediction remains unchanged for all allowed perturbations within a specified radius or norm bound.

checkpoint/restart A fault-tolerance strategy that periodically saves job state so a failed run can resume from the last saved state instead of starting over.

checkpoint storm A synchronized surge of massive network and storage write traffic that occurs when all nodes in a distributed training fleet flush model parameters and optimizer states simultaneously at periodic intervals.

chiplet A modular silicon die that is combined with other dies in one package (2.5D interposer or 3D stacking, such as TSMC CoWoS) so an accelerator can exceed the reticle limit of a single die.

chunked prefill A prompt ingestion strategy that divides long input contexts into fixed-size token chunks co-scheduled alongside single-token decode steps, preventing compute-intensive prefill passes from causing head-of-line latency spikes on decode streams.

circular economy An industrial economic model aimed at eliminating waste through the continual use, refurbishment, remanufacturing, and recycling of physical assets.

client drift The divergence of local model updates from the global optimal trajectory during federated learning, caused by multiple local optimization steps on heterogeneous, non-IID data distributions.

client scheduling The process of selecting which devices participate in federated learning rounds based on availability, data quality, and resource constraints to ensure representative model updates.

clock gating A circuit technique that suppresses the clock signal to inactive logic blocks, reducing dynamic power without changing the supply voltage or clock frequency of active blocks.

cloud ML Machine learning systems that use cloud computing infrastructure to provide scalable computational resources for training and inference, typically offering high-bandwidth connectivity and substantial processing power.

communication-computation ratio The ratio $\rho = T_{\text{comm}}(N)/(T_{\text{compute}}/N)$ of per-step communication time to local computation time. Values near or above 1 indicate substantial communication pressure; achievable overlap depends on independent work and available resources. The ratio is the time-domain form of communication intensity.

co-packaged optics (CPO) An advanced packaging architecture that co-locates silicon photonics transceivers directly on the same substrate as the network switch ASIC, shortening electrical trace lengths to millimeters to drastically lower interconnect energy per bit.

collective operation A group communication pattern in which all ranks participate to aggregate, broadcast, or redistribute data.

combinational logic Digital logic circuits where the output depends only on the current input states, not any past states or memory elements.

communication-computation overlap The concurrent execution of data transfer and useful computation whose dependencies are already satisfied. Only communication that fits within the available independent work can be hidden; any remaining transfer tail stays on the critical path.

communication intensity (CI) The ratio of bytes transferred across a network to local floating-point operations performed. It extends arithmetic-intensity reasoning from a device's memory hierarchy to a distributed system's network fabric.

compute express link (CXL) An open industry-standard interconnect protocol built on the PCIe physical layer that provides high-bandwidth, low-latency, cache-coherent shared memory access and pooling across heterogeneous processors and memory expanders.

concept bottleneck models Neural network architectures that first predict interpretable intermediate concepts before making final predictions, combining deep learning power with transparency.

concept drift A statistical change in the relationship between input features and target labels over time ($p(y \mid x)$), causing trained model predictions to degrade even if input data distributions remain unchanged.

confidential computing A security technology that protects data in use by performing computation inside hardware-isolated, cryptographically attestation-backed Trusted Execution Environments (TEEs).

conformal prediction A framework that uses held-out calibration data to construct prediction sets or intervals with finite-sample marginal coverage at a specified level under exchangeability.

congestion control Mechanisms that regulate traffic injection or routing in response to load, preventing persistent queues, packet loss, and tail-latency amplification in a shared fabric.

consistency The contract that determines which version a read may observe when writes overlap or replicas disagree.

consistent hashing A hash-based routing method that keeps keys or sessions mapped to stable replicas while minimizing remapping when replicas are added or removed.

constitutional AI A training method where models learn to improve their own outputs by critiquing responses against a set of principles, enabling iterative self-refinement and reducing harmful content while maintaining helpfulness.

contestability The architectural property and operational mechanism enabling end users or auditors to interrogate, challenge, and seek redress or recourse against algorithmic decisions.

context parallelism An attention-level distributed execution technique that partitions long sequence lengths across dedicated process groups using point-to-point ring communication or sequence-head transposes.

continual learning The ability of machine learning systems to learn continuously from a stream of data while retaining previously acquired knowledge, addressing the challenge of catastrophic forgetting in neural networks.

continuous batching An inference scheduling architecture that groups and schedules active requests at each autoregressive token generation iteration rather than waiting for an entire batch to complete, eliminating padding waste.

controller pattern An asynchronous, declarative control loop (standardized in Kubernetes) that continuously observes the actual state of cluster resources and drives mutations to converge toward the desired declared state.

cooling effectiveness The efficiency with which a data center cooling system removes heat from computing equipment, typically measured as the ratio of heat removed to energy consumed for cooling.

coordinated checkpointing A distributed fault-tolerance protocol where all participating worker ranks synchronize at a common execution step to capture a globally consistent snapshot of application state.

counterfactual explanation An interpretability mechanism specifying the minimal input perturbation required to flip a model's output prediction or decision.

covariate drift A shift in the distribution of input features ($p(x)$) between training and production environments while the true conditional probability distribution of targets ($p(y \mid x)$) remains unchanged.

covariate shift A type of distribution shift where the input distribution changes while the conditional relationship between inputs and outputs remains stable.

critical batch size A workload-dependent crossover beyond which increasing the batch yields diminishing reductions in the optimization steps needed to reach a target quality. The gradient noise scale can provide a proxy; the crossover is not necessarily a hardware-throughput or training-time optimum.

critical message size The operational crossover payload size ($n^* = \alpha \cdot \beta$) in the α-β model below which communication is dominated by startup latency and above which transfer time is dominated by link bandwidth.

CUDA graphs A GPU execution model that captures a sequence of kernel launches and memory dependencies into a static graph, amortizing CPU dispatch overhead through single-request replay.

D

D·A·M taxonomy A single-node systems framework decomposing workloads into Data (dataset volume and pipeline throughput), Algorithm (computational complexity and arithmetic intensity), and Machine (compute capabilities and memory hierarchy), projected onto distributed fleets via the C^3 taxonomy.

dark launching A deployment strategy where newly updated model versions process production traffic in the background without exposing outputs to end users, measuring system performance and stability under real load.

data center A facility housing computer systems and associated components such as telecommunications and storage systems, typically including redundant power supplies, cooling systems, and network connections.

data compression Techniques for reducing the size and complexity of training data through encoding, quantization, or feature extraction to enable efficient storage and processing on memory-constrained devices.

data contract A formal agreement between data producers and downstream consumers defining schema definitions, semantic constraints, and freshness expectations for ML feature pipelines.

data parallelism A training strategy where each worker holds a full model replica, processes a different data shard, and synchronizes gradients before applying the same update.

data poisoning An adversarial attack where malicious or corrupted data is deliberately injected into training datasets to degrade overall model accuracy or introduce specific vulnerabilities.

data provenance An immutable, auditable record tracking the origin, lineage, licensing, filtering criteria, and preprocessing transformations applied to training data throughout its lifecycle.

data sanitization The process of deliberately and permanently removing or destroying data stored on memory devices to make it unrecoverable, ensuring data security.

data stall ratio A diagnostic efficiency metric measuring the fraction of total training step duration during which accelerators sit idle waiting for input data from storage or preprocessing pipelines.

datasheets for datasets A standardized documentation framework that records the provenance, composition, collection methodology, and intended use cases of machine learning datasets.

decode phase The autoregressive generation phase in LLM inference that emits tokens one step at a time using the KV cache, often limited by memory bandwidth.

decode piggybacking A co-scheduling optimization in chunked prefill systems where memory-bandwidth-bound decode token generation executes concurrently with compute-bound prefill matrix multiplications, increasing GPU arithmetic intensity.

deep learning A subset of machine learning using neural networks with multiple hidden layers to automatically learn hierarchical representations from data.

defensive distillation A technique that trains a student model to mimic a teacher model's behavior using soft labels, reducing sensitivity to adversarial perturbations.

demographic parity A group fairness criterion requiring that the likelihood of receiving a positive algorithmic prediction or outcome is identical across all protected demographic subgroups.

Dennard scaling The historical observation that as transistors become smaller, their power density remains approximately constant, allowing for more transistors without proportional increases in power consumption.

differential privacy A mathematical guarantee bounding how much a randomized algorithm's output distribution can change between neighboring datasets. The parameters ϵ and δ limit individual influence under the stated adjacency relation; they do not promise that an observer learns nothing about a person.

differentially private SGD (DP-SGD) A training algorithm that clips per-example gradients and adds calibrated noise, with privacy accounting for sampling and repeated updates under a stated adjacency relation.

direct attach copper (DAC) A passive, shielded twinaxial copper cable assembly with fixed transceiver terminations used for short in-rack interconnections, introducing negligible transmission latency and zero active transceiver power consumption.

direct-to-chip liquid cooling A thermal management architecture that circulates liquid coolant through sealed cold plates mounted directly on high-power processors, capturing heat far more efficiently than air cooling.

disaggregated serving An inference serving architecture that physically separates prompt processing (prefill) and token generation (decode) onto dedicated, independently scaled hardware pools.

disparate impact ratio A quantitative fairness measure defined as the ratio of the positive selection rate for an underprivileged protected group to the positive selection rate for the privileged group, traditionally assessed against the four-fifths rule.

displacement of overhead The principle that reducing one distributed-systems cost often moves pressure to another part of the system rather than eliminating it. For example, fewer synchronization barriers can reduce coordination time while increasing staleness or recovery complexity.

distributed checkpointing A checkpointing pattern where workers write their own portions of training state in parallel under a coordinated protocol.

distributed tracing An observability technique that follows a request across multiple services or components using propagated trace context.

distributed training The practice of training machine learning models across multiple computing nodes or devices to reduce training time and enable larger model architectures.

distribution shift The phenomenon where data encountered during model deployment differs from the training distribution, potentially degrading model performance.

domain adaptation Machine learning techniques that enable models trained on one domain to perform well on a different but related domain, addressing distribution mismatch challenges.

dominant resource fairness (DRF) A multi-resource fair-sharing algorithm that equalizes the maximum share of any resource type (such as GPU, CPU, memory, or network bandwidth) allocated to competing users or workloads.

double binary tree An AllReduce algorithm implemented in NCCL that routes communication across two complementary spanning trees simultaneously, achieving $\mathcal{O}(\log N)$ latency while sustaining near-optimal bandwidth utilization.

dual modular redundancy (DMR) A fault-tolerance technique that duplicates a computation across two independent systems and compares their outputs. A disagreement detects a fault but cannot identify or correct the faulty result without another mechanism.

dual-use dilemma The challenge of mitigating misuse of technology that has both positive and negative potential applications, particularly relevant in AI security.

duck curve A daily electrical load profile graph showing a pronounced midday dip in net demand caused by peak solar energy generation, followed by a steep evening ramp when solar generation ceases.

durability The property that acknowledged data remains recoverable after the failures covered by the storage contract.

dynamic batching A serving technique that briefly aggregates incoming requests into batches to improve throughput while controlling added latency.

dynamic voltage and frequency scaling (DVFS) A power-management technique that adjusts a processor's supply voltage and clock frequency to trade execution speed for energy use and thermal headroom.

E

edge AI The deployment of artificial intelligence algorithms directly on edge devices like smartphones, IoT sensors, and embedded systems, enabling real-time processing without cloud connectivity.

edge computing A distributed computing paradigm that brings computation and data storage closer to data sources, enabling real-time processing with reduced latency and lower bandwidth requirements.

edge intelligence The paradigm of deploying machine learning inference, continuous adaptation, and distributed coordination directly on edge devices and near-edge systems under local power, memory, and connectivity constraints.

edge ML Machine learning systems that perform inference and sometimes training at the edge of networks, typically on resource-constrained devices like smartphones or embedded systems with limited computational power.

effective FLOP/s The useful throughput of a fleet after its peak FLOP/s is reduced by single-device efficiency (MFU), scaling efficiency, and runtime goodput, with each loss counted in exactly one factor.

efficiency frontier The Pareto-optimal trade-off curve across latency, throughput, cost, memory, and model accuracy beyond which no metric can improve without degrading another.

elastic training Training that can continue as workers join, leave, fail, or are replaced by resizing process groups and redistributing state or data.

elastic weight consolidation (EWC) A continual learning regularization algorithm that slows learning on parameters critical to previously mastered tasks by penalizing changes weighted by their Fisher information matrix values.

electromigration The movement of metal atoms in a conductor under the influence of an electric field, potentially causing permanent hardware faults over time.

elephant flow A persistent, high-throughput network flow, such as a large gradient AllReduce or checkpoint transfer, that sustains near line-rate transmission and creates structural bottleneck risks under hash-based routing.

embodied carbon The cumulative greenhouse gas emissions generated during the raw material extraction, manufacturing, transportation, and construction of hardware and physical data-center infrastructure.

energy efficiency The ratio of useful output to energy input, measuring how effectively a system converts energy into desired computational work.

energy roofline model An analytical model relating computational intensity and data movement to energy use under specified operation costs and accounting boundaries.

energy accounting boundary The requirement to evaluate useful computation against the total energy consumed by compute, cooling, and network movement. It complements the Fleet Law's time decomposition with an energy-side measure of distributed efficiency.

ensemble methods Machine learning approaches that combine predictions from multiple models to improve accuracy, robustness, and reliability compared to individual models.

environmental monitoring The systematic collection and analysis of environmental data using sensor networks and machine learning to track ecosystem health, pollution levels, and climate change impacts.

equal-cost multi-path (ECMP) A routing method that hashes flows across multiple paths with equal routing cost; persistent large flows can collide on one path even when other paths remain underused.

equal opportunity difference The difference in true positive rates between protected demographic groups for a binary classification decision.

equality of opportunity A fairness criterion focused on ensuring equal true positive rates across groups, guaranteeing that qualified individuals are treated equally regardless of group membership.

equalized odds A group fairness metric requiring that a predictor achieves identical true positive rates and identical false positive rates across all protected demographic groups.

erasure coding A mathematical data durability scheme that divides objects into d data and m parity fragments distributed across separate failure domains, providing extreme fault tolerance at lower capacity overhead than full replication.

error-correcting codes Methods used in data storage and transmission to detect and correct errors, improving system reliability and data integrity.

error feedback A distributed optimization technique for biased gradient compression (such as Top-k sparsification or 1-bit quantization) that stores compression residuals in local memory and adds them to subsequent updates, preserving convergence guarantees.

ESP32 A low-cost microcontroller unit widely used in IoT applications, featuring a 240 MHz processor and 520 KB SRAM, commonly deployed in resource-constrained social impact applications.

evaluation contract A specification defining quantitative performance thresholds, test suites, and regression checks that a machine learning model must satisfy before promotion to production.

event time The time at which an event occurred at its source, as distinct from the processing time at which a system handled it.

exactly-once outcome An application-level result in which replay cannot lose or duplicate an effect, typically enforced with stable identifiers and idempotent or transactional writes rather than transport delivery alone.

expected calibration error (ECE) A scalar metric that measures the discrepancy between predicted confidence probabilities and actual empirical accuracies across partitioned confidence bins.

experience replay A memory-based technique that stores past training examples in a buffer to prevent catastrophic forgetting and stabilize learning in streaming or continual adaptation scenarios.

expert collapse A training pathology in mixture of experts models where only a few experts receive significant training signal, causing other experts to become underutilized and reducing the model's effective capacity.

expert parallelism A distributed strategy for Mixture of Experts (MoE) architectures where different experts are assigned to different accelerators, requiring AllToAll collective routing to dispatch tokens to their assigned experts and combine results.

explainability The ability of stakeholders to understand how a machine learning model produces its outputs through post-hoc explanation techniques.

explicit congestion notification (ECN) A network-layer protocol extension that marks congestion indicators in IP packet headers when switch buffer queues exceed a preset threshold, triggering sender rate reductions without packet drops.

external memory Mechanisms that allow neural networks to access and manipulate external storage systems, extending their working memory beyond parameter storage to enable more complex reasoning and information retrieval.

F

fail-stop failure A failure mode where a computing component halts operation completely and visibly upon encountering a fault, ceasing all communication and computation.

failure in time (FIT) A reliability unit equal to one failure per 10^9 device-hours. A component's mean time to failure in hours is 10^9 divided by its FIT rate.

fairness constraints Technical and policy restrictions designed to ensure equitable treatment across demographic groups in machine learning systems.

fast gradient sign method (FGSM) A gradient-based adversarial attack that generates adversarial examples by adding small perturbations in the direction of the gradient.

fat-tree topology A Clos-style hierarchical network topology whose aggregate capacity increases toward the spine to provide high bisection bandwidth and multipath routing.

fault injection attack A physical attack that deliberately disrupts hardware operations through techniques like voltage manipulation or electromagnetic interference to induce computational errors and compromise system integrity.

fault tolerance The ability of a system to continue operating correctly even when some of its components fail or encounter errors.

feature store A centralized operational data management system that automates the ingestion, transformation, storage, and retrieval of curated features for both offline training and low-latency online inference.

federated averaging The standard algorithm for federated learning where client model updates are aggregated using weighted averaging based on local dataset sizes to produce a global model.

federated learning A machine learning approach that trains across decentralized data sources without centralizing raw training

records. Updates and released models can still reveal information, so privacy guarantees require additional mechanisms and explicit assumptions.

few-shot learning A machine learning paradigm that enables models to adapt to new tasks using only a small number of labeled examples, critical for data-sparse on-device scenarios.

FinOps A cloud financial management discipline and cultural practice that enables organizations to maximize business value by helping engineering, finance, and technology teams collaborate on data-driven spending decisions.

FlashAttention An exact, memory-efficient attention algorithm that computes softmax tiling in on-chip SRAM to eliminate quadratic intermediate writes to HBM.

Fleet Law The master equation of distributed step time, $T_{\text{step}}(N) = T_{\text{compute}}/N + T_{\text{comm}}(N) + T_{\text{sync}}(N) - T_{\text{overlap}}$, which decomposes an idealized fixed-work training step on N devices into the Compute, Communication, and Coordination terms of the C^3 taxonomy minus the time hidden by overlap. It is the distributed counterpart of the iron law.

fleet stack The book's layered model of a machine learning fleet, with eight rows (Facility & Accelerators, Fabric & Storage, Distributed Execution, Reliability & Control, Runtime & Serving, Operations, Assurance, and Governance & Sustainability) grouped into four bands (Infrastructure, Distribution, Serving, and Governance), one band per part of the book. Constraints at lower rows bound the choices available above them.

FLOP/s Floating-point operations per second, a throughput measure of how many floating-point arithmetic operations a system performs per second.

flowlet switching An adaptive routing scheme that divides network flows into discrete packet bursts (flowlets) separated by idle gaps longer than path delay skew, rerouting flowlets across alternative links without triggering packet reordering.

forward error correction (FEC) A digital channel-coding technique that transmits structured parity symbols alongside payload data, enabling the receiver to detect and correct bit transmission errors without requesting retransmissions.

foundation models Large-scale pretrained models like GPT and BERT that can be adapted for a wide variety of downstream tasks, serving as a foundation for multiple applications.

fully sharded data parallel (FSDP) A native PyTorch implementation of ZeRO-3 that shards model parameters, gradients, and optimizer states across ranks, using AllGather before layer computation and ReduceScatter during the backward pass.

G

gang scheduling A co-scheduling discipline that atomically allocates all required compute and network resources for a distributed job simultaneously, preventing hold-and-wait deadlocks in barrier synchronization.

GDPR The General Data Protection Regulation, a European Union law that imposes strict requirements on personal data processing and significantly influences privacy-preserving machine learning design.

generality tax The area, power, and control overhead a general-purpose processor spends supporting diverse workloads rather than accelerating the operations dominant in a specialized workload.

generative adversarial network (GAN) A model class in which a generator network learns to produce synthetic data while a discriminator network learns to tell it from real data; the two are trained against each other.

goodput The rate of useful, retained progress over a stated wall-clock interval and resource allocation. Training goodput can account for statistical efficiency, replay, and recovery, so high arithmetic utilization alone does not establish high goodput. A goodput claim names its measure, such as ML productivity goodput (MPG) for fleet capacity or SLO goodput for requests that meet their latency targets.

governance frameworks Structured approaches for managing responsible AI development including policies, procedures, oversight mechanisms, and accountability structures.

GPU Graphics Processing Unit, a specialized electronic circuit designed to rapidly manipulate and alter memory to accelerate the creation of images and parallel processing tasks like AI training.

GPUDirect RDMA A direct data path that lets RDMA-capable network adapters transfer data to or from GPU memory without staging through host memory.

GPUDirect Storage (GDS) An NVIDIA technology enabling a direct DMA data path between NVMe storage devices and GPU high-bandwidth memory, bypassing CPU host DRAM bounce buffers and reducing per-transfer latency.

graceful degradation The ability of a distributed system to continue operating in a reduced capacity rather than failing completely when one or more components become unavailable.

gradient accumulation A technique that accumulates gradients over multiple micro-batches before an optimizer step or synchronization.

gradient descent An optimization algorithm that iteratively updates model parameters by moving in the direction opposite to the gradient of the loss function, fundamental to neural network training.

gradient inversion attack An attack that reconstructs sensitive training examples or features from shared gradients or model updates.

gradient staleness The delay, measured in updates or steps, between when a gradient is computed and when it is applied.

graph break A fallback event in tracing JIT compilers where unsupported dynamic Python constructs force the compiler to terminate graph capture and return execution to the standard Python interpreter.

gray-box attack An adversarial attack where the attacker has partial knowledge about the model, such as knowing the architecture but not the specific parameters or training data.

gray failure A subtle, degraded operational failure where a component continues running and reporting healthy status while experiencing severe performance degradation or producing intermittent corruptions.

group distributionally robust optimization (GDRO) An optimization formulation that minimizes the worst-case training loss across predefined demographic or contextual subgroups rather than average empirical loss.

grouped-query attention (GQA) An attention mechanism where multiple query attention heads share a single key-value head, drastically diminishing the KV cache memory footprint per decoding step relative to Multi-Head Attention.

Gustafson's law The observation that when problem size grows with processor count, the scaled speedup grows nearly linearly, as $N - s(N-1)$ for serial fraction s of the scaled run. It describes weak scaling.

H

hallucination A generative model failure where output is fluent and confident but unsupported, fabricated, or factually wrong.

hardware redundancy The duplication of critical hardware components to provide backup functionality and improve system reliability through voting mechanisms.

hardware-software co-design The coordinated, simultaneous optimization of hardware microarchitecture, system software, and model algorithms to achieve performance and efficiency unattainable through independent component design.

hardware trojan A malicious modification embedded in hardware components during manufacturing that can remain dormant under normal conditions but trigger harmful behavior when specific conditions are met.

heartbeat A periodic liveness signal that a node or process sends to a monitor; a missed heartbeat is the usual trigger for failure detection and recovery.

hierarchical AllReduce A multi-phase collective reduction that exploits the cluster bandwidth gap by performing local reductions over fast intra-node interconnects before routing aggregated node shards across slower inter-node fabrics, finishing with intra-node result distribution.

high-bandwidth memory (HBM) A 3D-stacked DRAM architecture where vertically bonded memory dies connect to host processors via thousands of through-silicon vias (TSVs) across a silicon interposer, delivering terabytes-per-second memory bandwidth.

Hockney model Another name for the alpha-beta model of point-to-point communication, $T(n) = \alpha + n/\beta$, after the author who applied it to parallel computers.

homomorphic encryption A cryptographic technique that allows computations to be performed directly on encrypted data without decrypting it first, enabling privacy-preserving machine learning inference.

hot spares Backup components kept ready to instantaneously replace failing components without disrupting system operation, providing redundancy.

Huber loss A robust loss function used in regression that is less sensitive to outliers compared to squared error loss, improving training stability.

human-AI collaboration The synergistic partnership between humans and AI systems where each contributes their unique strengths to solve complex problems more effectively than either could alone.

human oversight The principle that human judgment should supervise, correct, or halt automated decisions, maintaining meaningful human control over AI systems.

hyperparameter optimization The process of finding the optimal configuration of hyperparameters (learning rate, batch size, network architecture parameters) that control the machine learning training process.

I

I/O wall The widening performance disparity where storage media bandwidth and host bus interfaces cannot keep pace with accelerator consumption rates, resulting in low device utilization and idle compute cycles.

idempotent write A write whose repeated execution with the same stable identifier produces the same stored result as one execution, making replay safe after uncertain failures.

immersion cooling A liquid cooling technique where entire server blades and electronic components are submerged directly in a thermally conductive, electrically dielectric fluid bath.

in-network reduction (SHARP) A hardware-assisted collective communication mechanism that offloads associative reduction operations directly to switch ASICs at line rate, eliminating intermediate host-memory round trips.

incast A many-to-one traffic burst in which simultaneous senders overwhelm one destination egress or receiver, rapidly filling shared buffers and increasing tail latency.

inference The phase in machine learning where a trained model is used to make predictions or generate outputs on new, previously unseen data.

InfiniBand A high-performance, low-latency interconnect widely used for HPC and ML training clusters, with native RDMA and credit-based flow control.

inherently interpretable model A model architecture whose internal decision logic is directly inspectable, transparent, and understandable by design without requiring post-hoc explanation techniques.

intermittent faults Hardware faults that occur sporadically and unpredictably, appearing and disappearing without consistent patterns, making diagnosis challenging.

interpretability The degree to which humans can understand the reasoning behind a machine learning model's predictions, often referring to inherently transparent models.

iron law The single-machine performance equation $T = D_{\text{vol}}/\text{BW} + O/(R_{\text{peak}}\eta_{\text{hw}}) + L_{\text{lat}}$, which divides execution time into data movement, computation, and fixed overhead. The Fleet Law is its distributed counterpart.

J

jailbreak attack A prompt manipulation technique designed to bypass safety filters, alignment guardrails, and system guidelines of large language models to elicit restricted or harmful outputs.

K

k-anonymity A privacy guarantee where each released record is indistinguishable from at least $k_{\text{anon}} - 1$ others after generalizing or suppressing quasi-identifiers.

knowledge distillation A technique where a large, complex "teacher" model transfers its learned knowledge to a smaller, more efficient "student" model, maintaining performance while reducing computational requirements.

Kueue A Kubernetes-native job queueing controller that manages multi-tenant job admission, resource quotas, and cross-cluster queue borrowing before pods are created on worker nodes.

KV cache A memory buffer that stores previously computed transformer key and value tensors so autoregressive generation does not recompute attention history for every token.

KV cache wall The operational bottleneck in large-scale inference where accelerator high-bandwidth memory capacity and interconnect transfer bandwidth become saturated by dynamic key-value states across long contexts, rather than by model parameter weights.

L

label shift A type of distribution shift in which the target-label distribution $p(y)$ changes while the class-conditional feature distribution $p(x \mid y)$ remains fixed.

large language models (LLMs) Neural networks with billions or trillions of parameters trained on vast text corpora, capable of understanding and generating human-like text across diverse domains and tasks.

life cycle assessment (LCA) A standardized, comprehensive cradle-to-grave accounting methodology (governed by ISO 14040/14044) for evaluating the environmental impacts of a product or service across design, manufacturing, operational use, and end-of-life disposal.

lighthouse archetype One of the three recurring workloads (Archetypes A, B, and C) that the book uses to expose a particular systems constraint across chapters. An archetype preserves continuity while allowing each chapter to examine the same workload through a different engineering lens.

Little's law The queuing identity $Q_{\text{req}} = \lambda_{\text{arr}} T_{\text{lat}}$, which relates the average number of requests in a stable system to their arrival rate and average time in the system.

load balancing Techniques in mixture of experts models to ensure that computational load and training signal are distributed evenly across experts, preventing expert collapse and maintaining model efficiency.

LogGP model An extension of the LogP communication model incorporating a gap-per-byte parameter (G) for large messages to accurately capture streaming and network pipelining effects in bulk data transfers.

LogP model An abstract parallel network performance model characterizing transfers using network latency (L_{LogGP}), processor injection overhead (o), message injection gap (g), and rank count (N).

lookup table A data structure that replaces runtime computation with simpler array indexing operations, commonly used for performance optimization.

LoRa A low-power, wide-area physical-layer radio modulation. LoRaWAN is the networking protocol that uses LoRa radios; achievable range depends on the frequency plan, transmit power, antennas, terrain, interference, and link budget.

low-rank adaptation A parameter-efficient fine-tuning method that approximates weight updates using low-rank matrices, reducing trainable parameters while maintaining adaptation capability.

M

machine learning A subset of artificial intelligence that enables computers to learn and make decisions from data without being explicitly programmed for each specific task.

machine learning fleet A distributed system of accelerators, storage, network fabrics, and control software engineered to operate as one coherent computational machine.

machine learning lifecycle The complete process of developing, deploying, and maintaining ML systems, from data collection through model retirement.

machine learning operations (MLOps) The practice of deploying and maintaining machine learning models in production reliably and efficiently through automated pipelines.

machine learning security The protection of data, models, and infrastructure from unauthorized access, manipulation, or disruption throughout the entire machine learning lifecycle.

machine unlearning Methods for removing or reducing the influence of specified training records from a model. The claimed guarantee must be evaluated against a stated reference training procedure; an approximate update does not automatically establish complete removal.

marginal emissions The emissions caused by an incremental increase in electricity demand, often from the generator dispatched next on the grid.

mean time between failures (MTBF) The average operating time between consecutive failures of a repairable component or system. For N independent components with constant failure rates, system MTBF is the component value divided by N.

mean time to failure (MTTF) The average operating time before a component fails. Under a constant failure rate λ, it equals $1/\lambda$.

mean time to repair (MTTR) The average time required to diagnose, repair or replace, and restore a failed system component to full operational status.

megawatt-hour A unit of energy equal to one megawatt of power used for one hour, commonly used to measure electricity consumption in large facilities like data centers.

membership inference attack An adversarial privacy attack where an adversary queries a trained model to determine whether a specific individual's record was part of the model's training dataset.

meta-learning The process of learning how to learn, where models are trained to quickly adapt to new tasks with minimal data, particularly useful for personalization in on-device systems.

mice flow A transient, low-volume network flow, such as an RPC barrier or heartbeat message, whose primary performance constraint is tail latency and queuing delay rather than raw bandwidth.

microcontroller A compact integrated circuit designed to govern specific operations in embedded systems, typically featuring limited processing power and memory but optimized for low power consumption and real-time applications.

minimax A decision-making strategy used in game theory that attempts to minimize the maximum possible loss in adversarial scenarios.

mixed-precision training A technique that uses different numerical precisions for different parts of neural network training, typically combining 16-bit and 32-bit floating-point arithmetic to reduce memory usage and increase training speed.

mixture of experts An architectural approach that uses multiple specialized sub-models (experts) with a gating mechanism to route inputs to the most relevant experts, enabling efficient scaling while maintaining sparsity.

ML productivity goodput (MPG) The fraction of a fleet's capacity-time that becomes useful work at the hardware's ideal rate, the product of scheduling, runtime, and program goodput. Each factor is measured on a fixed accounting boundary so that every lost interval is charged to exactly one factor.

mobile ML Machine learning systems optimized for mobile devices like smartphones and tablets, balancing computational efficiency with inference accuracy for on-device processing.

MobileNet A family of efficient neural network architectures designed for mobile devices using depthwise separable convolutions to achieve significant reductions in model size and computation.

mode collapse A failure mode in generative models where the model produces only a limited variety of outputs, ignoring the diversity present in the training data and failing to capture the full distribution.

model card Standardized, short documentation reports that disclose the intended usage, performance benchmarks across subgroups, limitations, and ethical considerations of machine learning models.

model compression Techniques to reduce the size and computational requirements of machine learning models through methods like quantization, pruning, and knowledge distillation to enable deployment on resource-constrained devices.

model extraction An attack where an adversary reconstructs or clones the proprietary architecture, weights, or capabilities of a target model by systematically querying its public API.

model FLOPs utilization (MFU) The fraction of hardware peak FLOP/s spent on useful model computation over a stated measurement window. At fleet scale it measures program goodput only, how close retained steps run to the hardware peak, not the efficiency of the fleet as a whole.

model inversion attack A privacy attack that reconstructs representative features or private training data samples by analyzing model outputs and confidence scores.

model quantization A technique that reduces the precision of model parameters (typically from 32-bit to 8-bit or lower) to decrease model size and computational requirements while maintaining acceptable accuracy.

model registry A central catalog for model artifacts, versions, metadata, lineage, dependencies, and lifecycle state.

model theft An attack that copies a deployed model's parameters, architecture, or behavior, either exactly (from leaked weights or artifacts) or approximately (by training a substitute on query responses, also called model extraction).

model uncertainty The inadequacy of a machine learning model to capture the full complexity of the underlying data-generating process, leading to prediction uncertainty.

model watermarking A technique for embedding verifiable ownership signatures into machine learning models that can be used to detect unauthorized use or prove intellectual property theft.

Monte Carlo dropout A technique that uses multiple forward passes with different dropout masks at inference time to estimate prediction uncertainty.

Moore's law The observation that the number of transistors on a microchip doubles approximately every two years, reducing cost per transistor under historical scaling assumptions.

multi-instance GPU (MIG) NVIDIA's hardware-level partitioning architecture that slices a physical GPU into up to seven fully isolated GPU instances, each with dedicated memory, SMs, and crossbar bandwidth.

multi-process service (MPS) NVIDIA's runtime multiplexing layer that allows multiple CUDA processes to execute concurrently on a single GPU by sharing execution engines, compute contexts, and memory bandwidth.

multicalibration A fairness technique ensuring that model predictions remain calibrated across intersecting subgroups, addressing complex demographic interactions.

N

NCCL NVIDIA Collective Communications Library, a GPU communication library for topology-aware collective operations.

network fabric The links, switches, routing, and transport mechanisms that connect compute nodes and determine the bandwidth, latency, and failure behavior of distributed communication.

neural architecture search Automated methods for designing optimal neural network architectures, using algorithms to explore the space of possible network designs and find the best performing structures.

neural engine Specialized hardware accelerators designed for machine learning inference and training, such as Apple's Neural Engine or Google's Edge TPU, optimized for on-device AI workloads.

neural network A computational model inspired by biological neural networks, consisting of interconnected nodes (neurons) organized in layers that can learn complex patterns from data.

neural processing unit (NPU) A dedicated domain-specific hardware accelerator optimized for tensor arithmetic, low-precision operations, and high energy efficiency in mobile and edge devices.

neuromorphic computing Computing architectures inspired by the structure and function of biological neural networks, designed to process information more efficiently than traditional digital computers.

non-blocking fabric A network whose internal capacity allows any feasible permutation of endpoint pairs to communicate at full line rate without contention inside the fabric.

non-IID data Training or evaluation data across distributed devices whose local distributions are not independent and identically distributed, inducing client drift and slowing convergence in federated learning.

Non-Volatile Memory Express (NVMe) A high-performance protocol for SSDs attached over PCIe, often used as a local cache or checkpoint staging tier in ML storage hierarchies.

NVLink A high-bandwidth NVIDIA interconnect for GPU-to-GPU communication within a node.

NVMe over Fabrics (NVMe-oF) A storage networking protocol that extends NVMe command sets and queue parallelism across low-latency RDMA or TCP fabrics, enabling remote cluster storage access at near-local solid-state latency.

O

object storage A storage model that manages data as objects with metadata and network API access, commonly used for durable lower-cost capacity.

on-device adaptation The local training or fine-tuning of a subset of parameters of a pre-trained model directly on edge silicon without transmitting private training examples off-device.

on-device learning The local adaptation or training of machine learning models directly on deployed hardware devices without reliance on continuous connectivity to centralized servers.

online softmax A numerically stable, incremental softmax algorithm that updates running maximum and sum statistics across tiles, eliminating the need to materialize global attention matrices.

operational carbon The greenhouse gas emissions produced directly and indirectly by the consumption of electrical energy to run computing systems and facility infrastructure during their operational lifespan.

operator fusion An optimization that combines adjacent operations into fewer kernels so intermediate values stay on chip and memory traffic or launch overhead falls.

optical circuit switching (OCS) A physical networking architecture that redirects optical signals directly between fiber ports using micro-electromechanical mirror (MEMS) arrays without intermediate optical-electrical-optical conversions.

orchestration The coordination of fleet resources and workloads, including scheduling, placement, scaling, sequencing, recovery, and resource allocation across distributed infrastructure.

out-of-distribution (OOD) detection The algorithmic task of identifying whether a given input sample originates from a different data distribution than the dataset on which the model was trained.

over-selection A federated client scheduling strategy where the coordinator selects more clients than required for a training round and aggregates the first responding cohort to mitigate stragglers.

oversubscription A topology design in which aggregate downstream capacity exceeds upstream capacity, reducing switch cost while creating a shared bandwidth bottleneck under simultaneous traffic.

oxide breakdown The failure of an oxide layer in transistors due to excessive electric field stress, causing permanent hardware faults.

P

packet spraying A dynamic multipath routing technique that disperses individual packets belonging to the same transport flow across multiple equal-cost network paths simultaneously, maximizing bisection bandwidth utilization.

PagedAttention An attention computation algorithm and dynamic memory manager inspired by operating system virtual memory paging that partitions key-value cache tensors into non-contiguous physical page blocks linked via page tables.

PAM4 Pulse Amplitude Modulation 4-Level, a multilevel signaling scheme that encodes two bits per baud across four distinct electrical voltage levels, doubling channel throughput over traditional NRZ encoding.

parallel file system (PFS) A distributed file system (such as Lustre or GPFS) that decouples metadata operations from striped data blocks across multiple storage nodes to aggregate high concurrent bandwidth.

parameter A learnable variable in a machine learning model, such as weights and biases

in neural networks, that are adjusted during training to optimize performance.

parameter-efficient fine-tuning (PEFT) Adaptation methods that update only a small subset of parameters or added modules while keeping most pretrained weights frozen.

partition skew An uneven distribution of records or keys that leaves one distributed-pipeline partition with disproportionate work and makes it determine stage completion time.

permanent faults Hardware defects that persist irreversibly until repair or component replacement, consistently affecting system behavior.

personalization layers Model components, typically the final classification layers, that are adapted locally to user-specific data while keeping shared backbone layers frozen.

PFC deadlock A permanent buffer starvation condition in a multi-hop lossless Ethernet network where cyclical priority flow control PAUSE-frame dependencies among switch queues halt packet forwarding indefinitely.

PFC storm A cascading congestion phenomenon in lossless fabrics where buffer saturation at a single slow-draining endpoint causes priority flow control PAUSE frames to propagate upstream across multiple tiers, freezing unrelated victim flows.

physical attack Direct manipulation or tampering with computing hardware to compromise the security and integrity of machine learning systems, bypassing traditional software defenses.

pipeline bubble Idle accelerator time in pipeline parallelism while stages wait for the pipeline to fill or drain. For N_{PP} balanced stages and m microbatches, the bubble fraction is $(N_{\text{PP}}-1)/(N_{\text{PP}}-1+m)$.

pipeline parallelism A distributed training strategy that partitions layers into sequential stages and uses microbatches to overlap work across devices.

placement group A scheduler abstraction that enforces spatial or topological co-location constraints across multiple workers, minimizing cross-switch communication latency in distributed model training.

pluggable optics Modular, field-replaceable optical transceiver units (such as QSFP and OSFP form factors) mounted to switch and NIC faceplates that interface SerDes electrical traces with optical fibers.

point-in-time correctness The requirement in feature computation that training datasets reflect the exact feature values available at the precise historical timestamp of the prediction event, preventing data leakage.

point-in-time join A join that retrieves feature values valid for a historical prediction and available at that prediction time, preventing future or late-arriving information from leaking into training.

post-hoc explanations Explanation methods applied after model training that treat the model as a black box and infer reasoning patterns from input-output behavior.

power delivery network (PDN) The physical electrical distribution hierarchy, spanning rack-scale busbars, board power planes, voltage regulator modules, and package decoupling capacitors, that supplies stable sub-volt, kilo-ampere current to low-voltage silicon gates.

power usage effectiveness (PUE) The ratio of total data-center facility energy to the energy delivered to IT equipment. A value of 1.0 would mean no energy spent on cooling, power conversion, or other facility overhead.

predictive parity A fairness criterion requiring that the positive predictive value of a classifier is equal across different protected demographic subgroups.

preemption cascade A systemic failure mode where preempting one job causes cascading resource evictions across dependent or lower-priority jobs, incurring compounding checkpointing and cold-start overheads.

preemption overhead The cumulative unproductive overhead incurred by preempting and restarting a distributed training job, comprising time spent checkpointing, reloading weights, reconstructing communication groups, and recomputing lost iterations.

prefetch buffer A host DRAM queue that stages and holds pre-decoded, augmented training batches ahead of accelerator execution to absorb storage access jitter and tail latency spikes.

prefill-decode disaggregation A distributed serving architecture that physically decouples prompt prefill processing onto compute-dense GPU nodes from token decoding onto memory-bandwidth-optimized GPU nodes, transferring intermediate KV cache activations across network fabrics.

prefill phase The initial LLM inference phase that processes the prompt and populates the KV cache, often dominated by compute.

prefix caching A KV cache optimization that detects and retains precomputed attention activations for shared prompt prefixes (such as system instructions or few-shot exemplars), bypassing redundant prefill computation across sessions.

principle of least privilege A security concept where users are given the minimum access levels necessary to complete their job functions, reducing security risks.

priority flow control (PFC) An IEEE 802.1Qbb link-layer flow-control standard for lossless Ethernet that pauses upstream senders via PAUSE frames on a designated traffic class before switch ingress buffers saturate.

privacy budget The differential-privacy parameters governing the permitted change in output distributions on neighboring inputs. The parameter ϵ controls a multiplicative bound of e^{ϵ}, δ allows additive slack, and privacy accounting tracks how guarantees compose across releases.

privacy-preserving techniques Methods intended to protect privacy at specified boundaries, including differential privacy and cryptographic protocols. Federated learning and local processing reduce raw-data movement but do not alone guarantee that updates or outputs reveal no private information.

privacy-utility trade-off The fundamental tension between preserving individual privacy and maintaining the utility of data for machine learning, requiring careful balance through techniques like differential privacy.

program goodput The factor of ML productivity goodput (MPG) that measures how close retained steps run to the hardware's ideal rate, charging exposed communication, pipeline bubbles, stragglers, and memory stalls. MFU measures it for compute-bound training.

projected gradient descent (PGD) An iterative gradient-based adversarial attack that repeatedly perturbs an input and projects it back into an allowed norm ball.

prompt injection An attack against large language models where untrusted input content overrides system instructions and directs the model to perform unauthorized actions.

pruning A model compression technique that removes unnecessary connections or neurons from neural networks to reduce model size and computational requirements without significantly impacting performance.

Q

quantization The process of reducing the precision of model weights and activations from floating-point to lower-bit representations to decrease memory usage and accelerate computation on edge devices.

quantization-aware training (QAT) Training that simulates low-precision arithmetic so a model remains accurate after deployment with quantized weights or activations.

quota governance The organizational policies and control mechanisms (such as guaranteed vs. burst allocations, review cycles, and utilization-based reclamation) that regulate multi-tenant capacity reservation.

R

rail-optimized routing A communication mapping that aligns corresponding model shards with the same local GPU and network rail across nodes. It can distribute traffic across available NICs, but shared links, placement, and competing traffic still determine contention.

rail-optimized topology A cluster network topology tailored for multi-accelerator servers where identical GPU ranks across nodes are connected to separate, dedicated leaf-switch fabrics, minimizing switch traversals for data-parallel collectives.

randomized smoothing A certified-defense method that classifies noisy versions of an input to obtain probabilistic robustness guarantees.

RDMA over Converged Ethernet (RoCE) A protocol that carries remote direct memory access semantics over Ethernet, relying on carefully managed low-loss or lossless behavior.

recursive halving-doubling A collective reduction algorithm that pairs ranks at exponentially increasing distances across $\log_2 N$ rounds, achieving both logarithmic latency and bandwidth optimality for power-of-two participant counts.

red teaming A structured adversarial testing exercise where security and safety experts deliberately probe ML systems, interfaces, and serving pipelines for safety vulnerabilities, security flaws, and compliance failures.

ReduceScatter A collective communication operation that reduces inputs across ranks and scatters distinct shards of the reduced result back to ranks.

regularization Methods used in machine learning to prevent overfitting by adding penalty terms to the loss function, constraining model complexity.

reliability collapse The exponential decline, $A_{\text{component}}^{N}$, in the probability that all N required components of a tightly coupled job are available, even when each component's availability $A_{\text{component}}$ is high.

reliability gap The difference between component reliability and the much lower probability that every component required by a large, tightly coupled job remains available. The gap widens as fleet size and job duration increase.

remote direct memory access (RDMA) A networking mechanism that lets one machine

read or write another machine's memory without kernel-mediated copies on the data path.

renewable energy Energy collected from renewable resources that are naturally replenished, including solar, wind, hydroelectric, geothermal, and biomass sources.

responsible AI The practice of developing and deploying AI systems in ways that are ethical, fair, transparent, and beneficial to society while minimizing potential harms and biases.

reward hacking The phenomenon where AI systems exploit unintended aspects of reward functions to maximize scores while violating the intended objectives.

ridge point The arithmetic intensity in the Roofline Model at which the memory bandwidth ceiling meets the peak computational ceiling, delineating the boundary between memory-bound and compute-bound operational regimes.

ring AllReduce A bandwidth-optimal collective communication algorithm that organizes participants into a logical ring, executing sequential ReduceScatter and AllGather neighbor transfers so each node communicates a constant volume of data independent of rank count.

RLHF Reinforcement learning from human feedback, a training method that uses human preferences to guide model behavior, enabling AI systems to better align with human values and intentions.

robust AI The ability of artificial intelligence systems to maintain performance and reliability despite internal errors, external perturbations, and environmental changes.

robustness A model's ability to maintain stable and consistent performance under input variations, environmental changes, or adversarial conditions.

robustness margin The minimum distance or perturbation magnitude required to alter a model's prediction on a given input sample under defined perturbation norms.

robustness metrics Quantitative measures for evaluating model stability under various perturbations, including adversarial accuracy, certified robustness bounds, and performance under distribution shift.

roofline model A performance model that compares arithmetic intensity with peak compute and memory bandwidth to identify compute-bound and bandwidth-bound regimes.

runtime goodput The factor of ML productivity goodput (MPG) that measures the share of held time that becomes retained progress at the job's fault-free step rate, charging startup, input stalls, checkpoint pauses, failure recovery, rework, and preemption.

S

scale moment A critical threshold where an ML workload outgrows the memory capacity, arithmetic throughput, or reliability envelope of a single node, necessitating distributed execution and exposing systems to the network wall, energy wall, and reliability collapse.

scaling efficiency The ratio of actual distributed training throughput to ideal linear speedup ($\eta_{\text{scaling}} = T_1/(N \cdot T_N)$), quantifying capacity loss from communication, synchronization barriers, and idle pipeline bubbles.

scaling laws Empirical power-law relationships between model loss and training compute, parameter count, or dataset size, of the form $\mathcal{L}(X) = A_{\mathcal{L}} X^{-\alpha_{\text{scale}}} + C_{\mathcal{L}}$, valid within the regime in which they were fitted. They explain why each further gain in model quality requires multiplicatively more resources.

scheduling goodput The factor of ML productivity goodput (MPG) that measures the share of fleet capacity-time held by complete allocations, charging unallocated capacity, partially assembled gangs, and fragmentation.

scope 1 emissions Direct greenhouse gas emissions from sources that are owned or controlled by an organization, such as on-site diesel backup generators and company facilities.

scope 2 emissions Indirect greenhouse gas emissions associated with the generation of purchased electricity, steam, heating, or cooling consumed by an organization.

scope 3 emissions All other indirect greenhouse gas emissions that occur throughout an organization's upstream and downstream value chain, including hardware manufacturing and supply chain logistics.

secure aggregation A cryptographic protocol that lets a coordinator recover an aggregate of participating clients' updates without seeing individual unmasked updates, under stated participation and collusion assumptions. It does not by itself bound information revealed by the aggregate or final model.

secure boot A hardware-backed startup process that cryptographically verifies firmware and software before execution.

secure computation Cryptographic protocols that enable multiple parties to jointly compute functions over private inputs without revealing those inputs to each other.

secure multi-party computation A cryptographic method that allows multiple parties to jointly compute a function over their private inputs without revealing those inputs to each other.

self-supervised learning A machine learning paradigm where models learn representations from unlabeled data by predicting parts of the input from other parts, reducing dependence on manually labeled datasets.

sequence parallelism A distributed technique that shards activations along the token sequence dimension across tensor-parallel ranks during non-GEMM operations (LayerNorm, Dropout) using ReduceScatter and AllGather without increasing total communication volume.

service-level objective (SLO) A measurable target for service behavior such as latency, freshness, availability, or model quality.

shadow deployment A deployment mechanism where production traffic is duplicated and routed to a candidate model alongside the active production model to validate performance and latency without affecting user experience.

SHAP SHapley Additive exPlanations, a feature-attribution framework based on Shapley values from cooperative game theory.

sharded checkpointing A checkpointing approach where each worker saves only its assigned model, gradient, or optimizer shard.

sharded data parallelism A distributed training strategy (exemplified by ZeRO and FSDP) that partitions model parameters, gradients, and optimizer states across data-parallel ranks, reconstructing full layers on demand via collective communication.

side-channel attack An attack that exploits information leaked through the physical implementation of computing systems, such as power consumption, electromagnetic emissions, or timing variations.

silent data corruption (SDC) A hardware error in compute units, memory, or data paths that changes a stored or computed value without raising any detectable error, so the corrupted result propagates silently into training or inference.

single instruction, multiple threads (SIMT) A parallel execution model where multiple independent threads execute the same instruction stream in lockstep across vector execution units, hiding high memory latency through hardware warp scheduling over large thread pools.

small file problem A data-loading bottleneck where accessing millions of individual sample files exhausts storage metadata server request capacity long before physical data transfer bandwidth saturates.

software fault Unintended behavior in software systems resulting from defects, bugs, or design oversights that can impair performance or compromise security.

sparse training A training approach that maintains sparsity in neural network weights throughout the training process, reducing computational requirements and memory usage.

sparse updates A training strategy that selectively updates only a subset of model parameters based on their importance or contribution to performance, reducing computational and memory overhead.

specification gaming When AI systems find unexpected ways to achieve high rewards that technically satisfy the objective function but violate the intended purpose.

speculative decoding An inference acceleration mechanism where a compact, lightweight draft model speculatively generates multiple candidate tokens in parallel that are subsequently verified in a single forward pass by the large target model.

speculative execution A performance optimization in processors that executes instructions before confirming they are needed, which can inadvertently expose sensitive data through microarchitectural side channels.

split learning A distributed learning paradigm and model partitioning technique where a neural network is divided between an edge client and a central server, with intermediate activations and gradients exchanged at a cut layer.

stale synchronous parallel (SSP) A synchronization model that lets faster workers advance a bounded number of steps ahead of the slowest worker.

straggler A correct but slow worker that delays synchronous progress for the rest of a distributed job.

strong scaling Adding processors to a fixed problem to finish it sooner. Its speedup is bounded by Amdahl's law.

stuck-at fault A permanent hardware fault where a signal line becomes fixed at a logical 0 or 1 regardless of input, causing incorrect computations.

supply chain attack An attack that compromises hardware or software components during the manufacturing, distribution, or integration process, potentially affecting multiple downstream systems.

sustainable AI The practice of developing and deploying artificial intelligence systems that minimize environmental impact while maintaining effectiveness and accessibility.

switch radix The number of ports on a network switch, which constrains endpoint fan-out, topology depth, and the number of switching tiers needed at a given cluster scale.

symbol rate The number of signaling symbols transmitted per second on a link; multilevel modulation carries multiple bits per symbol without increasing this physical transition rate proportionally.

synchronous tight coupling A distributed execution dependency where all participating nodes must complete local computation and synchronize states at global barriers before any worker can proceed, coupling cluster progress to the slowest device.

synthetic data generation The creation of artificial datasets that approximate the statistical properties of real data while reducing privacy risks and avoiding direct exposure of sensitive information.

systolic array A specialized grid of tightly coupled processing elements that rhythmically computes and streams matrix data directly between neighbor cells without writing intermediate partial sums back to register files or memory.

T

tail latency High-percentile request latency, such as p99, that captures rare slow responses affecting distributed service reliability.

targeted attack A type of data poisoning attack that aims to cause misclassification of specific inputs or classes while leaving the model's general performance largely intact.

tensor core A specialized accelerator execution unit designed to perform mixed-precision fused matrix multiply-accumulate operations ($\mathbf{D} = \mathbf{A} \cdot \mathbf{B} + \mathbf{C}$) on matrix tiles as a single hardware instruction.

tensor parallelism An intra-layer model parallelism technique that shards individual weight matrices and operations (such as column- and row-parallel linear layers in multi-head attention and MLPs) across accelerators connected via high-bandwidth interconnects.

thermal design power (TDP) The vendor-specified maximum sustained rate of heat energy in watts that a cooling subsystem must dissipate under maximum expected operating workloads to maintain silicon below its critical junction temperature.

threat modeling A structured security analysis that identifies assets, adversary capabilities, attack paths, and defensive priorities for a system.

tiered storage A hierarchical data placement architecture spanning hot local NVMe, warm parallel file systems, cold object storage, and deep archive tiers to optimize throughput against capacity economics.

time per output token (TPOT) The average wall-clock time required to generate each subsequent token during autoregressive decoding, governed primarily by accelerator memory bandwidth.

time to first token (TTFT) The latency elapsed between submitting an inference request and receiving its initial generated token, reflecting the duration of the compute-bound prefill phase and serving as a primary interactive responsiveness metric.

time-travel query A database or feature store capability that reconstructs the exact state of feature data as it existed at any historical point in time.

TinyML Machine learning architectures, runtimes, and optimization techniques engineered to execute inference and ultra-lightweight adaptation on microcontrollers operating under milliwatt power and sub-megabyte memory limits.

topology-aware scheduling Placement that accounts for GPU and interconnect locality such as NVLink domains, racks, rails, or switch hierarchy.

total cost of ownership (TCO) A comprehensive financial assessment of the complete direct and indirect costs associated with designing, deploying, and maintaining machine learning infrastructure throughout its operational lifecycle.

TPU Tensor Processing Unit, a specialized AI accelerator chip developed by Google specifically designed for machine learning workloads, particularly neural network computations.

trackable resource (TRES) A fine-grained accounting and scheduling abstraction (standardized in Slurm) that treats individual accelerators, licenses, memory, and interconnect rails as distinct consumable scheduling units.

training The process of teaching a machine learning model to make predictions by showing it examples and adjusting its parameters based on performance feedback.

training-serving skew A discrepancy in performance between training and production environments caused by differences in data processing pipelines, feature definitions, or serving runtime environments.

transfer learning A machine learning technique that uses knowledge from a pretrained model on one task to improve learning on a related task, enabling efficient adaptation with limited data.

transformer A neural network architecture that uses self-attention mechanisms to process sequential data, forming the foundation for many modern language models.

transient faults Temporary hardware faults that do not persist or cause permanent damage but can lead to incorrect computations if not handled properly.

transient voltage droop A rapid drop in core supply voltage caused by the product of parasitic loop inductance and high current slew rates ($\Delta V = L_{\text{parasitic}} \frac{di}{dt}$) when thousands of execution cores abruptly transition between idle barriers and dense computation.

transparency Openness about how AI systems are built, trained, validated, and deployed, including disclosure of data sources, design assumptions, and limitations.

tree AllReduce A latency-oriented collective algorithm that reduces data up a spanning tree to a root before broadcasting results down, achieving $\mathcal{O}(\log N)$ round-trip latency at the expense of higher link bandwidth consumption on large payloads.

triple modular redundancy (TMR) A fault-tolerance technique that runs three replicas of a computation and uses majority voting to mask one independent faulty result; it does not guarantee correctness under multiple or common-mode faults.

Triton A domain-specific language and compiler for tiled neural network programming that abstracts thread scheduling and memory coalescing while exposing block-level memory hierarchy management.

two-phase commit (2PC) A coordination protocol with prepare and commit phases that ensures participants commit atomically or abort.

U

uncoordinated checkpointing A checkpointing scheme where distributed nodes record local state independently without cross-node synchronization, risking domino-effect rollbacks during recovery.

V

value alignment The principle that AI systems should pursue goals consistent with human intent and ethical norms, addressing the challenge of encoding human values in machine objectives.

value-sensitive design A methodology for incorporating human values into technology design through systematic stakeholder engagement and ethical consideration of system impacts.

Volcano A cloud-native batch scheduling system for Kubernetes designed for distributed ML and big-data workloads, providing gang scheduling, fair-share queuing, and resource reservation primitives.

W

wafer-scale engine (WSE) An accelerator architecture that integrates compute, memory, and communication across most or all of a silicon wafer to reduce the bandwidth and latency costs of crossing separate package boundaries.

wake-word detection An always-on keyword spotting workload executed on low-power edge processors that detects specific acoustic trigger phrases with minimal latency and sub-milliwatt power draw.

warehouse-scale computer (WSC) A building-scale computing system operated as one coherent machine, with the network fabric and storage hierarchy treated as system-level resources.

watchdog timer A hardware component that monitors system execution and triggers recovery actions if the system becomes unresponsive or stuck.

water usage effectiveness (WUE) The liters of water a data center consumes per kilowatt-hour of IT energy, determined largely by its cooling and heat-rejection design.

watermark A streaming-system estimate that event time has advanced far enough to close a window, subject to an explicit policy for arrivals that appear later.

weak scaling Growing the problem in proportion to the number of processors so that work per processor stays roughly constant. Gustafson's law describes its speedup.

WebDataset A high-throughput sequential data loading format that packages dataset samples inside standard POSIX tar archives, enabling direct linear streaming over object storage without per-file filesystem overhead.

weight freezing A technique that fixes most model parameters during training while allowing only specific layers or components to be updated, reducing computational requirements for on-device adaptation.

white-box attack An adversarial attack where the attacker has complete knowledge of the model's architecture, parameters, training data, and internal workings, enabling highly effective attack strategies.

Y

Young-Daly formula The first-order optimal checkpoint interval $T^*_{\text{ckpt}} = \sqrt{2 T_{\text{write}} \, \text{MTBF}_{\text{system}}}$, which balances checkpoint write time against expected

rework after failures. It is accurate when writes are short compared with the system MTBF.

Z

ZeRO (Zero Redundancy Optimizer) A memory-optimization paradigm that progressively eliminates redundant state across data-parallel workers by partitioning optimizer states (ZeRO-1), gradients (ZeRO-2), and model parameters (ZeRO-3).

zero-bubble pipeline scheduling An advanced pipeline scheduling strategy that decouples backward passes into activation gradient computation (B) and deferred weight gradient computation (W), using the W kernels to fill startup, turnaround, and drain bubbles.

zero-day vulnerability A previously unknown security flaw in software or hardware that can be exploited by attackers before developers have had a chance to create and distribute a patch.

Index

1-bit Adam, 247
  variance freeze, 247
1-bit SGD, 244
1F1B scheduling
  memory efficiency, 297
  peak activation bound, 297
  steady-state execution, 297
2.5D interposer
  CoWoS, 100
24/7 carbon-free energy
  definition, 949
3D integration
  bandwidth density, 1041
3D parallelism
  accelerator allocation, 263
  collective primitives, 215
  communication intensity reduction, 13
  definition, 263
  grid decomposition, 309
  traffic patterns, 130
  training loop, 311
48 V DC distribution
  distribution losses, 925

A/B testing
  power, 705
  statistical power, 705
A100, 260
Abstention
  safety mechanism, 1002
Accelerator
  selection, 96
  topology, 1032
Access pattern
  sequential vs. random, 154
Accountability
  governance mechanisms, 966
Activation checkpointing
  appendix refresher, 1082
  edge memory, 616
  full recomputation, 268
  memory footprint, 78
  memory reduction, 268
  recomputation overhead, 268
  selective recomputation, 268
  traffic cost, 622
Activation outliers
  quantization degradation, 563
Active optical cable, *see* AOC
Adam
  distributed update, 269
Adam optimizer
  checkpoint memory, 348
  checkpoint state size, 187
  state, 348
Adaptation strategy
  trade-offs, 634
Adapter switching
  definition, 630
Adaptive batching
  request-aware grouping, 541
  SLO-adaptive feedback, 541
  traffic-adaptive window, 541
Adaptive ensemble
  concept drift response, 866
Adaptive evaluation
  gradient masking pitfalls, 878
Adaptive routing
  queue-depth selection, 139
Admission control
  definition, 394
  edge adaptation, 623
  overload shedding, 581
Advanced packaging
  reticle limit, 100
  stacking yield, 60
Adversarial attack
  autonomous driving stop sign, 872
  definition, 869
  input-space manipulation, 849
  prompt evasion, 824
Adversarial defense
  four-layer workflow, 874
Adversarial example
  discovery, 785
Adversarial patch
  physical road sign attack, 889
  physical-world attack, 872
Adversarial perturbation
  decision boundary geometry, 869
Adversarial training
  accuracy cost, 875
  definition, 874
  free gradient reuse, 874
  PGD batch augmentation, 874
Adversarial transferability
  black-box, 801
AES
  AES-256, 829
  side-channel vulnerability, 810
Aggregation
  small-file mitigation, 168
Aggressive sharing
  definition, 434
AI compute growth
  doubling time, 897
AI hardware e-waste
  definition, 897
Air-cooling limit, 1105
Air-gapped
  isolation, 779
Aleatoric uncertainty
  distinction from epistemic, 877
Alert fatigue
  clinical workflow, 1012
  monitoring thresholds, 720
AlexNet
  two-GPU training, 5
Algorithm
  D·A·M taxonomy, 25
Algorithm aversion
  human-AI interaction, 1013
Algorithmic bias
  historical data, 971
  proxy variable selection, 972
  reproduction in training, 971
Algorithmic fairness
  definition, 971
Algorithmic stability
  large-batch dynamics, 262
Alignment regression testing
  guardrail validation, 1004
Alignment tax
  capability trade-off, 1003
All-pairs bandwidth matrix
  definition, 144
AllGather
  decomposition, 218
  definition, 218
  fabric traffic pattern, 109
  FSDP parameter prefetch, 277
  parameter reconstruction, 268
  ring topology, 223
Allocation shape
  definition, 405
AllReduce, 280
  algorithm selection, 255
  asynchronous execution, 250
  bandwidth lower bound, 223
  convergence guarantee, 213
  data-parallel gradient averaging, 266
  decentralized gradient synchronization, 268
  decentralized scaling, 263
  definition, 218
  etymology, 212
  fabric traffic pattern, 109
  gradient exchange volume, 262
  hierarchical, 269
  inter-node, 235
  overlap with backward pass, 251
  recursive doubling, 227
  ring algorithm, 223, 280
  scaling ceiling, 272
  tensor parallelism synchronization, 567
  tensor-parallel decode, 494
AllToAll
  bipartite exchange, 241
  combine phase, 241
  definition, 218
  dispatch phase, 241
  embedding communication, 60
  embedding lookup, 308
  expert parallelism, 305, 571
  fabric traffic pattern, 109
  hierarchical exchange, 242
  injection bound, 239
  pairwise exchange schedule, 243
  reception bound, 239
  routing matrix, 241
  schedule selection, 242
  variable-size AllToAllv, 241
Alpha-beta model
  crossover message size, 1114
  link bandwidth, 1113
  message size, 1113
  message transfer time, 121
  startup latency, 1113
Amazon S3
  key-value namespace, 171
Amdahl's law, 281
  fixed residual fraction, 281
  parallel scaling limit, 101
Amdahl, Gene, *see also* Amdahl's law
Anderson-Darling test, 877
Anomaly detection
  percentage change, 756

percentage change threshold, 756
Z-score, 756
Z-score threshold, 756
AOC
definition, 116
Apache Arrow
in-memory format, 164
Apache Parquet
columnar storage, 164
API, 213
API keys
credential leakage, 822
Apple Neural Engine
definition, 619
on-device training, 619
Application monitoring
definition, 143
Application programming interface, *see* API
Application-specific integrated circuit, *see* ASIC
Approximate behavioral theft
definition, 786
Archetype
A (GPT-4/Llama-3), 9
A (GPT-4/Llama-3), compute infrastructure, 60
B (DLRM at Scale), compute infrastructure, 60
B (DLRM at Scale), pod design, 95
C (Federated MobileNet), compute infrastructure, 57
Archetype B
DLRM scaling, 307
goodput decomposition, 502
Architectural debt
delayed governance, 1042
Architecture
disaggregated, 147
Arduino
memory constraint, 617
Arduino Nano 33 BLE Sense
definition, 617
Arithmetic intensity
algorithmic, 488
appendix refresher, 1079
decode phase, 560
definition, 58
FLOP-to-byte ratio, 63
measured, 488
performance engineering, 488
single-device vs fleet, 263
vs. communication intensity, 24
ARM TrustZone
adoption, 816
ASIC, 237
ASP, *see* Asynchronous SGD
Assumptions
capacity planning, 1153
cluster tiers, 1147
nameplate capacity, 1155
network, 1147
reliability, 1149
sustainability, 1151
Assurance lifecycle
evidence requirements, 963
Asynchronous checkpointing, *see* Checkpointing
Asynchronous collective
overlap limits, 255
Asynchronous coordination
federated learning, 653
Asynchronous execution
non-blocking communication, 494
Asynchronous SGD
staleness penalty, 270
Attack
deployment-time, 782
inference-time, 782
postdeployment, 782
training-time, 782
Attained service scheduling
queueing discipline, 429
two-dimensional, 429
Attention
maps in security monitoring, 824
Audit trail
decision justification, 1026
user-level provenance, 993
Auditability
definition, 25
Automation
automation cost, 682
Automation bias
human-AI interaction, 1013
AutoML
hyperparameter optimization, 761
Automotive safety integrity level
ML certification, 852
Autopilot
perception failure, 851
Autoregressive bottleneck
memory bandwidth bound, 533
Autoregressive generation
variable sequence length, 533
Autoscaling
capacity control, 585
hysteresis, 586
predictive scaling, 587
Auxiliary load balancing loss
router skew, 306
Auxiliary loss, 571
Availability
definition, 24
MTBF and MTTR, 328
storage, 171
Average absolute odds difference
audit metric, 978
AWQ
activation-aware scaling, 564
salient weight protection, 564
AWS S3 outage
cascade failure, 851
AWS Trainium
accelerator fabric, 49

Backdoor attack
BadNets, 797
trigger pattern, 883
trigger patterns, 797
Backfill scheduling
conservative backfilling, 401
reservation, 401
Backpressure
data pipelines, 193
propagation, 581
Bad placement
definition, 416
Balls-into-bins
queue length bound, 576
Bandwidth
bound, 220
effective, 235
lower bound, 223
network, 280
operating regime, 221
optimality, 222
per-node PFS, 189
storage hierarchy, 156
Bandwidth hierarchy
distance scaling, 70
Bandwidth-delay product
buffer distinction, 140
Bandwidth-distance product
optical reach, 213
Bandwidth-dominated communication, 1114
Barrier synchronization
slowest worker dependence, 10
Batch size
computation-to-communication ratio, 279
Batched gather-matrix-vector multiplication
definition, 555
Batching
diminishing returns, 1139
efficiency curve, 528
knee of the curve, 528
latency knee, 1139
max batch size, 541
padding waste, 533
preferred batch size, 541
stability threshold, 1139
timeout parameter, 541
utilization headroom, 1139
Bathtub curve
infant mortality, 338
useful life, 338
wear-out, 338
Bayesian neural network
uncertainty estimation, 877
Benchmarking
end-to-end measurement, 497
placement context, 256
BERT
gradient communication, 281
TPU cluster training, 5
BF16
dynamic range, 67
mixed-precision training, 67
BGMV
definition, 555
kernel launch, 556
Bias
amplification feedback loop, 1012
automation, human-AI interaction, 1013
Bias-only adaptation
definition, 626
Bin packing
multi-dimensional, 398
scheduling, 398
Binding constraint
cross-layer analysis, 1043
Biofilm
definition, 91
Bisection bandwidth
AllToAll communication, 240
definition, 124
in communication axis, 26
network fabric, 1035
network wall, 4
scaling ceiling, 1034
Block quantization, 244
block scale factor, 244
Block table
logical to physical mapping, 547
Block-wise quantization
definition, 562
storage overhead, 563
Blue-green deployment
blue-green exposure, 710
ML serving, 687
Bonferroni correction
multiple comparison control, 863
Bottleneck
displacement, 596
Boundary attack
decision-based black-box, 872
Breakdown point
definition, 804
Broadcast, 218
tree topology, 218
BSP
synchronization, 270
Bucket fusion, 251
size tuning, 251
trade-off, 251
Budget
burn rate, 756
Buffered asynchronous aggregation
mechanism, 653
Build vs. buy
storage economics, 197
Build-vs.-buy decision
utilization threshold, 740
Bulk synchronous parallel, *see* BSP
definition, 111
origin, 111
Bulkhead pattern
isolation, 584
Burn rate

definition, 722
Burn-in testing
reliability screening, 338
Busbar
48V DC distribution, 86
Butterfly algorithm, 227
Byte-pair encoding
subword tokenization, 15
Byzantine failure
classification, 333
Byzantine fault tolerance
federated aggregation, 667
gradient aggregation, 852
Krum algorithm, 334
ML overhead, 334
trimmed mean, 334
Byzantine-resilient ML
algorithmic overhead, 334

Cache
effective service time, 158
Cache eviction
least recently used (LRU), 554
Calibration
expected calibration error, 867
fairness trade-off, 972
group-specific, 971
score, 974
selective prediction, 867
temperature scaling, 867
Canary deployment
canary, exposure, 710
global rollout, 591
gradual rollout, 711
risk mitigation, 761
safety validation, 1007
CAP theorem
definition, 24
scheduling, 393
Capacity factor, 571
expert routing, 241
mixture of experts, 306
Capacity planning, 531
bottom-up sizing, 745
earlier availability, 747
infrastructure requirements, 745
peak load vs. average load, 602
CapEx
cluster breakdown, 739
facility capex, 741
hardware, 741
network, 741
Capital expenditure
definition, 739
Carbon accounting
attributional, 903
consequential, 903
Carbon Border Adjustment Mechanism
definition, 954
Carbon budget
absolute limits, 943
fleet constraint, 1037
Carbon emissions
grid intensity, 20
Carbon footprint
benchmarks, 929
tracking, 911
training emissions, 896
Carbon intensity
by energy source, 915
regional grid variance, 926
Carbon intensity variance
definition, 912
Carbon usage effectiveness
data-center metric, 927
Carbon-aware scheduling
definition, 950
fleet placement, 1037
grid intensity APIs, 951
load shifting, 951
marginal vs. average emissions, 951
results, 951
Carlini-Wagner attack
defense evaluation, 871
optimization formulation, 871
Catastrophic forgetting
definition, 636
Causal effect
path-specific, fairness, 977
total, fairness, 977
CCPA
data deletion, 990
Certified robustness
definition, 876
limits of post-hoc filtering, 849
mathematical guarantee, 876
Channel pipelining, 253
Chaos engineering
serving resilience, 383
Checkpoint consistency
in-flight gradients, 385
partial writes, 385
rank desynchronization, 385
Checkpoint storm
definition, 188
fan-out read comparison, 155
write bursts, 155
Checkpoint/restart, 1099
Checkpointing
async PFS copy, 189
asynchronous, 351
automated recovery, 362
bounded asynchronous, 354
cadence and recovery, 1035
consistency cut, 354
coordinated, 352
definition, 347
eventual consistency, 355
format optimization, 191
in-memory staging, 351
incremental, 368
local checkpoint staging, 163
manifest publication, 190
overhead, 189
overhead trade-off, 1035
restore validation, 191
retention policy, 191
sharded, 355
sharded I/O parallelism, 355
shared namespace coordination, 168
storage architecture, 187
strict synchronous, 354
synchronous, 351
training pause, 189
uncoordinated, 352
Young-Daly formula, 1038
zero-redundancy, 191
Checksum
object storage verification, 174
Chi-square goodness-of-fit test
definition, 863
Chiplet
architecture, 51
die-to-die interconnect, 51
dual-die packaging, 51
manufacturing yield, 51
Chunked prefill
definition, 557
CI/CD
continuous integration, 697
patterns by model type, 718
Circuit breaker
circuit breaker pattern, 581
safety mechanism, 1007
serving resilience, 376
Circular economy
definition, 942
extended support cycles, 943
hardware lifespan extension, 942
modular design, 942
right-to-repair, 942
secondary-use programs, 943
Clean-label poisoning
definition, 798
label audits, 879
low-dose attacks, 889
public data sources, 889
Client drift
federated learning, 642
Client scheduling
federated learning, 644
Clock gating
definition, 84
Clos network, *see* Fat-tree
Closed-loop governance
feedback loop dynamics, 1012
Cloud computing
ML infrastructure, 744
Cluster
geometry, 117
MTBF scaling, 330
Cluster design
gradient bus integration, 141
Cluster management
pre-flight checks, 719
Cluster orchestration
distributed training, 263
Cluster utilization
communication bottleneck, 443
data starvation, 443
memory fragmentation, 443
zombie jobs, 443
Clustered federated learning
client grouping, 649
CMOS power equation
digital circuit foundation, 903
Co-packaged optics
architecture, 146
bandwidth hierarchy flattening, 147
link budget effects, 146
power savings, 146
thermal sensitivity, 147
CodeCarbon
definition, 954
Coefficient of variation
definition, 397
Cold restart
recovery trade-offs, 360
Cold start
LLM replica loading, 587
recommendation system, 761
Cold-start problem
definition, 657
personalization, 657
Collective communication
AllToAll schedules, 242
Collective operation
definition, 217
Columnar format
feature selection, 164
Communication efficiency
accuracy per byte, 663
client participation rate, 648
gradient compression, 646
Communication intensity
appendix refresher, 1108
definition, 13
fleet scaling, 262
Communication overhead
gradient synchronization, 262
interconnect bandwidth, 573
Communication overlap
collective hiding, 1035
Communication wall
gradient bus, 109
Communication-bound regime
efficiency drop, 282
Communication-bound workload, 1092, 1108
Communication-computation overlap
boundary effects, 252
definition, 249
energy efficiency, 282
fabric role, 108
incast mitigation, 141
latency hiding, 141, 271
LogP analysis, 249
memory overhead, 252
Communication-computation ratio
definition, 262
COMPAS
algorithmic bias, 995
Compilation cache
persistence across launches, 509
Compiler

cost model limitations, 514
Compiler-derived partitioning
definition, 490
Computational storage
SSD offload, 200
Compute demand
growth versus efficiency, 894
Compute Express Link
memory expansion, 200
Compute-bound, 488
Compute-limited regime
resource allocation, 18
Compute-optimal training
resource allocation, 15
Concept drift
definition, 858
detection algorithms, 859
fleet monitoring, 724
ground-truth feedback, 859
monitoring, 1041
remediation cost, 858
Confidence calibration
drift detection, 661
Confidence rounding
definition, 790
Confidential computing
hardware isolation, 805
Conformal prediction
exchangeability requirement, 868
finite-sample coverage, 867
guaranteed coverage, 868
routing and fallback, 890
Congestion control
PFC and ECN counters, 145
Congestion spreading
definition, 135
Consistency
definition, 24
read-after-write, 174
storage, 171
Consistent hashing
cache locality, 578
KV cache, 578
stateful routing, 578
Constitutional AI
alignment, 1003
Container isolation
ML serving, 825
Contestability
system accountability, 999
Contestability stack
appeal routing, 1001
architecture, 1001
decision provenance, 1001
explanation generation, 1001
outcome tracking, 1001
Context eviction
KV cache pressure, 582
Context filtering
safety mechanism, 1003
Context parallelism
AllToAll transpose, 304
definition, 302
Continual learning
edge devices, 656
Continuous batching
definition, 533
iteration-level scheduling, 98, 533
slot reuse, 533
Continuous integration
ML systems, 697
Contrastive learning
definition, 655
Contribution analysis
layer selection, 632
Control plane
definition, 25
fleet management, 672
Controller pattern
Kubernetes, 407
Coolant distribution unit
cooling loop isolation, 88
N+1 redundancy, 91
Cooling
capacity constraint, 1041
cooling class, 87
immersion, 88
junction temperature, 85
rack-scale heat dissipation, 81
Coordinated checkpointing, *see* Checkpointing
Coordination overhead
collective latency, 282
Copy-on-write
prefix caching, 548
Corporate Sustainability Reporting Directive
definition, 954
value chain reporting, 954
Correlated failure
cooling system, 334
failure domain protection, 334
network switch, 334
operator error, 334
power supply, 334
software bug, 334
Coscheduling
Kubernetes scheduling, 407
Cost
annual cost, 741
anomaly root cause categories, 756
per active user, 756
quality trade-off analysis, 757
Cost allocation
equal split, 755
marginal cost, 755
proportional, 755
Cost efficiency
scheduling objective, 395
Cost per inference
definition, 755
Cost policy
definition, 428
Counterfactual explanation
model interpretability, 997
recourse, 997
Counterfactual fairness
definition, 976
observational limitations, 977
Counterfeit hardware
provenance, 813
Covariate drift
importance weighting, 857
Covariate shift
day-night illumination, 857
definition, 857
fleet monitoring, 724
importance weighting, 857
CPO, *see* Co-packaged optics
CPU offloading
optimizer state, 78
CRC
gradient validation, 340
Credit-based flow control
definition, 119
Critical batch size
data-parallel ceiling, 272
definition, 286
scaling regimes, 286
Critical materials
AI hardware, 940
definition, 940
Critical message size, 220
algorithm crossover, 231
crossover, 254
Crossbar switch, 73
NVSwitch, 74
Crossover arithmetic intensity
compute and memory balance, 908
Crowdsourcing
poisoning risk, 796
Cryptographic hash
validation cost, 193
CSRD, *see* Corporate Sustainability Reporting Directive
CUDA, 849
programming model, 48
streams, 494
CUDA Graphs
appendix refresher, 1081
fleet gain, 496
CUDA time-slicing
definition, 433
CUE, *see* Carbon usage effectiveness
Curriculum learning
data locality, 155
Custom accelerator
cloud options, 744
CXL, *see* Compute Express Link
coherent interconnect, 100
definition, 100
memory pooling, 100
Cybersecurity regulations
ML compliance, 813
Cyclic buffer dependency
definition, 135
C^3 taxonomy
appendix refresher, 1085
binding constraints, 1039
communication, 26, 1092
compound loss, 1097
compute, 26, 1092
coordination, 26, 1093
definition, 26
operational dimensions, 25
pairwise intersections, 27
projection from D·A·M, 27
workload archetypes, 1039

DAC
definition, 116
signal attenuation loss, 117
DALI, *see* NVIDIA Data Loading Library
Daly formula, *see* Young-Daly formula
D·A·M taxonomy
single-node foundation, 1032
Dark launching
shadow traffic, 761
Data
data freshness, 696
limited regime, 18
Data augmentation
techniques, 882
workspace memory, 161
Data center
construction critical path, 746
environmental justice, 971
industrial electricity demand, 897
Tier III concurrent maintainability, 92
Tier IV fault tolerance, 92
Data center emissions scale
definition, 923
Data center industrial scale
definition, 897
Data contract
continuous validation, 697
definition, 696
Data drift, 858
incident classification, 729
Data efficiency
edge constraints, 624
scaling saturation mitigation, 20
Data format
compression integration, 165
index trade-offs, 164
Data gravity
placement economics, 183
Data leakage
feature store risk, 693
Data lineage
poisoning rollback, 882
Data loader
multi-worker throughput, 182
object storage streaming, 173
pinned memory, 161
pipeline architecture, 161
worker count, 182
Data loading
pipeline, 80
Data locality
privacy feature, 609
Data movement
energy cost, 60, 281

Data Nutrition Project
  dataset documentation, 1022
Data parallelism, 215
  batch loop transformation, 263
  communication cost, 323
  definition, 265
  inter-node, 71
  network hierarchy, 141
  pure, 77
  rail-optimized topology, 129
  synchronous SGD, 213
  trade-offs, 1038
Data pipeline
  throughput equation, 175
Data placement
  dynamic on-demand, 183
  static allocation, 183
Data poisoning
  availability attack, 797, 883
  backdoor attack, 797
  bilevel optimization, 796
  concept poisoning, 799
  definition, 796
  efficiency, 785
  online feedback manipulation, 879
  preview, 25
  subpopulation attack, 797
  support-vector machines, 879
  targeted attack, 797
  targeted class compromise, 883
  training integrity, 796
Data provenance
  ML pipelines, 882
  regulatory compliance, 174
Data quality
  production validation, 695
Data shuffling
  quality vs. I/O efficiency, 182
  shard-level, 182
Data staging
  on-demand staging, 183
  prestaging, 183
  streaming, 183
Data stall ratio
  definition, 176
  diagnostic metric, 177
Data transfer cost
  cloud egress, 196
Data wall
  definition, 195
Dataset scale
  storage limit, 260
Dataset versioning
  cryptographic snapshotting, 193
Datasheets for datasets
  operating envelope, 1009
  training data documentation, 1009
DCQCN
  congestion control, 138
  definition, 138
DDP
  comparison with FSDP, 277
  gradient bucketing, 263
Deadlock
  Coffman conditions, 401
  hold-and-wait, 401
  pipeline parallelism, 344
Debug port
  JTAG vulnerability, 812
Decentralized SGD
  neighbor communication, 291
Declarative model
  definition, 406
Decode
  inference phase, 98
Decode bottleneck, 1036
Decode phase
  memory bandwidth bottleneck, 557
  memory bandwidth bound, 527
  memory bandwidth saturation, 602
Decode piggybacking
  definition, 560
  weight loading amortization, 560
Deep ensemble
  uncertainty estimation, 877
Deep Gradient Compression
  gradient sparsification, 290
Deep leakage from gradients, 827
DeepFool
  minimal perturbation, 871
DeepSpeed
  1-bit Adam, 247
  Ulysses context parallelism, 304
DeepSpeed Ulysses, *see* DeepSpeed
Defense in depth
  definition, 782
  layered architecture, 815
  ML lifecycle integration, 855
Defense selection framework
  deployment contexts, 841
Defensive distillation
  definition, 877
Defragmentation
  definition, 399
Deionized water
  definition, 91
Delta encoding
  dataset snapshot, 202
Demographic parity
  definition, 972
  final prediction, 972
Denial of wallet
  availability abuse, 802
Dennard scaling
  power wall, 83
  sustainability, 898
  voltage limits, 83
Deployment
  automation payback period, 682
  consistency models for multi-region, 715
  dark launching, 761
  patterns by scale, 687
  risk mitigation, 710
  risk-based rollout, 712
  shadow mode, 702
  staged ensemble, 686
Deployment tier
  cloud ML, 967
  edge ML, 967
  mobile ML, 967
  TinyML, 967
Depreciation
  accelerator hardware, 743
Determinism
  violation, 333
Device mesh
  compiler-derived partitioning, 491
Device plugin
  definition, 407
DGX
  H100 baseboard, 73
Die-to-die interconnect
  bandwidth, 51
Differential privacy
  advanced composition, 835
  central, 828
  composition with secure aggregation, 651
  data sensitivity distribution, 839
  definition, 831
  edge devices, 650
  effective sample size, 839
  federated learning, 609
  Gaussian mechanism, 834
  Laplace mechanism, 833
  local, 828
  mean-estimation noise, 834
  moments accountant, 836
  noise injection, 650
  privacy accounting log, 651
  privacy budget epsilon, 832
  privacy loss random variable, 836
  Rényi differential privacy, 835
  simple composition, 835
  task complexity, 839
Dimension-ordered reduction, 236
Diminishing returns, 1142
Direct attach copper, *see* DAC
Direct preference optimization
  memory budget, 316
Direct-to-chip liquid cooling
  cold plate mechanism, 88
DiskANN
  SSD-backed vector index, 200
Disparate impact ratio
  threshold analysis, 977
Displacement of overhead
  cross-layer optimization, 1042
  cross-layer trade-offs, 1039
  definition, 29
  parallelism trade-offs, 324
DistBelief
  distributed training, 263
Distributed Checkpoint
  manifest publication, 354
  rank-sharded data, 191
Distributed ML
  governing principles, 1032
Distributed ML engineering
  fleet stack shift, 1032
Distributed scheduling
  challenges, 446
Distributed tracing
  context propagation, 379
  ML pipelines, 726
  observability, 379
  span, 379
  trace, 379
Distributed training
  definition, 261
DistributedDataParallel, 250, *see* DDP
  gradient hooks, 250
Distribution band
  fleet stack, 260
Distribution shift
  clinical imaging, 857
  definition, 857
  sensor noise, 848
District heating
  definition, 927
  municipal heat networks, 928
DLRM
  AllToAll, 215
  coordination bottleneck, 1039
  embedding scale, 307
  goodput decomposition, 502
  lighthouse archetype, 33
  scale, 520
DMR
  fleet contrast, 342
Domain shift, 857
Dominant resource fairness
  multi-resource allocation, 395
  progressive filling, 398
  worked allocation, 397
Double binary tree, 228
  complementary trees, 228
DP-SGD
  lot, 837
  training cost, 992
  training overhead, 832
DPO, *see* Direct preference optimization
DPU
  infrastructure offload, 75
Dragonfly
  topology origin, 131
Dragonfly topology
  definition, 131
DRAM
  data pipeline staging, 161
  energy cost, 907
DRF, *see* Dominant resource fairness
Drift detection
  on-device models, 661
  sample requirements, 725
Dropout
  MC, inference cost, 877
Dual-store architecture
  definition, 689
Dual-use dilemma
  ML security, 799

Duck curve
  definition, 951
  training pause during peak ramp, 952
Durability
  eleven nines, 171
  storage, 171
Duty cycle
  edge inference power, 911
  vs. goodput, 513
DVFS
  accelerator power management, 84
Dynamic batching, 527
  sequence length variance, 317
Dynamic power
  definition, 904
Dynamic random-access memory, *see* DRAM
Dynamic scaling
  FP8 dynamic range, 291
Dynamic sharing
  definition, 435
D·A·M taxonomy
  algorithm, 25
  data, 25
  definition, 25
  machine, 25
  workload projection, 27

E4M3
  weights and activations, 291
E5M2
  scaled gradients, 291
ECC
  end-to-end limitation, 386
  error detection and correction, 339
ECC error
  uncorrectable, 361
ECMP
  elephant flow limitations, 112
  hash collision, 139
  hash polarization, 112
  static flow hashing, 139
ECN
  CNP feedback loop, 138
  congestion signaling, 119
  marking threshold, 138
  reaction window, 140
  tuning sensitivity, 138
Edge AI
  power budget categories, 934
Edge AI energy trade-off
  definition, 913
Edge caching
  central inference, 591
Edge computing
  energy trade-off, 913
  selective hardening, 850
Edge deployment
  scaled-down cloud fallacy, 36
Edge inference
  vs. on-device learning, 608
Edge intelligence
  definition, 608
Edge learning
  integration controls, 657
  paradigm, 639
Edge power measurement
  instrument comparison, 910
EDP, *see* Energy-delay product
Effective cost
  definition, 425
Effective FLOP/s, 1097
Effective service rate, 1136
Effective training time
  Llama 3 window, 460
Effective training time ratio
  failure-driven composition, 481
Efficiency
  distributed, 31
  energy dimension, 31
  facility, 905
  power, 744
Efficiency frontier
  definition, 510
  Pareto optimality, 510
Egress fee
  cloud storage, 196
Elastic net attack to DNNs
  definition, 871
  sparse perturbation, 871
Elastic scheduling
  definition, 419
  elastic range, 418
  transition cost, 418
Elastic training
  definition, 364
  dynamic membership, 1035
  fleet fault tolerance, 1100
  resilience, 387
Elastic weight consolidation
  catastrophic forgetting, 866
  definition, 656
Electricity
  annual electricity, 741
Electromigration
  wear-out mechanism, 337
Electronic waste
  AI hardware obsolescence, 941
Elephant and mice flows
  definition, 111
Elephant flow
  ECMP collisions, 139
  gradient AllReduce, 111
  network traffic, 111
Embedded AI
  definition, 942
Embedding sharding
  column-wise, 307
  hot keys, 572
  hybrid, 307
  lookup deadline, 572
  parameter server, 307
  row-wise, 307
Embedding table
  host DRAM capacity, 161
  hybrid sharding, 597
  recommendation scale, 572
  sparse feature lookup, 539
Embodied carbon
  amortization formula, 918
  hardware lifecycle, 913, 1037
  hardware manufacturing, 913
Emissions trading
  compute cap, 954
Emissions trading for compute
  definition, 954
Emissions trading systems
  definition, 954
Empirical scaling law
  resource requirements, 14
Encrypted inference
  homomorphic encryption, 828
  latency overhead, 829
Energy
  fleet consumption metric, 31
  privacy trade-off, 968
Energy accounting boundary
  definition, 31
Energy buffering
  definition, 900
  transient smoothing, 926
Energy consumption
  data movement hierarchy, 214
Energy harvesting
  power budgets, 936
Energy reuse effectiveness
  heat export, 928
Energy roofline model
  crossover arithmetic intensity, 908
  definition, 908
Energy wall
  thermodynamic limit, 4
Energy-delay product
  energy-latency trade-off, 948
Energy-to-communication ratio
  local vs. cloud offloading, 935
Engineering practices
  fleet scale, 32
Enhanced transmission selection
  definition, 143
Ensemble
  price sensitivity model, 688
  product embedding model, 688
  ranking model, 688
  staged deployment, 686
  user embedding model, 688
Environmental justice in data center siting
  definition, 955
Epistemic uncertainty
  robustness, 877
Equal opportunity difference
  audit metric, 978
Equal-cost multi-path, *see* ECMP
Equality of opportunity
  true-positive rate, 973
Equalized odds
  error-rate parity, 972
Erasure coding
  definition, 172
  Reed-Solomon codes, 171
Error correction code, *see* ECC
Error feedback
  definition, 245
  residual accumulator, 245
  telescoping sum, 246
Error model
  propagation, 382
ESP32
  definition, 619
EU AI Act
  auditability requirement, 25
  energy reporting, 954
  high-impact FLOP threshold, 954
  risk-based classification, 965
  transparency and oversight, 965
EUV lithography energy cost
  definition, 913
Evaluation contract
  definition, 716
  production release, 1039
Event time
  data pipelines, 194
Event-driven computing
  energy comparison, 947
Exact artifact theft
  definition, 786
Exclusive access
  definition, 434
Executive Order 14110
  safety reporting thresholds, 966
Experience replay
  device tiers, 657
Expert buffering, 571
Expert collapse
  routing failure, 571
Expert parallelism, 570
  conditional computation, 304
  definition, 305
Explainability
  system evidence, 966
Explicit congestion notification, *see* ECN

Fabric behavior
  definition, 134
Fail-safe defaults
  definition, 782
Fail-stop failure
  characteristics, 333
  detection and recovery, 333
Failsafe mechanism
  ML serving, 852
Failure
  inevitability at scale, 387
Failure detection
  heartbeat protocol, 356
Failure domain
  atomic unit of failure, 383
  coupling patterns, 10
  definition, 80
  ML infrastructure, 334
  physical hierarchy, 80
  top-of-rack switch, 126
Failure modes

gray failure, 719
Failure rate
FITs, 327
fleet scale, 8
Failure-domain-aware placement
replica distribution, 170
Fair-share scheduling
decay half-life, 405
definition, 405
Fairness
accuracy cost, 972
gerrymandering, 979
impossibility theorems, 975
incompatibility, 971
release gate, 701
scheduling objective, 395
stability constraint, 966
Fairness impossibility theorem, 1036
Fake profile injection
definition, 880
sybil network, 880
Fallback redundancy
definition, 852
edge constraints, 852
ensemble switching, 854
toxicity guards, 851
Fallback routing
graceful degradation, 582
Fast gradient sign method
adversarial training, 870
definition, 870
single-step formulation, 870
Fat-tree
bisection bandwidth, 127
Clos origin, 127
definition, 126
pod fabric, 93
Fault diagnosis
cross-layer correlation, 887
fixed-input replay, 887
Fault injection
abstraction gap, 382
testing discipline, 382
tools and frameworks, 382
Fault injection attack
activation manipulation, 809
Fault model
abstraction, 382
Fault tolerance
cluster MTBF, 262
definition, 326
detection speed, 383
resource overhead, 850
spectrum, 384
FBLearner Flow
DAG orchestration, 761
Feature interaction, 539
recommendation systems, 539
Feature squeezing
adversarial defense, 877
bit-depth reduction, 877
definition, 877
Feature store
definition, 672
feature complexity, 753
feature staleness, 729
freshness requirements, 690
point-in-time correctness, 193, 693
recommendation systems, 524
training-serving skew, 689
Feature-parallel batching, 539
FEC
definition, 115
optical latency budget, 116
FedAvg
coordination cycle, 641
definition, 640
Federated A/B testing
client cohort validation, 661
Federated analytics
edge fleet, 662
privacy-preserving telemetry, 661
Federated coordination
edge constraints, 624
Federated learning, 840
compute-to-data paradigm, 609
definition, 609
edge constraints, 320
fairness constraint, 984
fleet heterogeneity, 609
introduction, 33
FedProx
proximal regularization, 642
Fidelity
fault injection framework, 382
Field-programmable gate array, *see* FPGA
Filter bubble
recommendation system, 762
Fine-tuning
selective, on-device, 632
FinOps
definition, 754
FIPS 140-2
security levels, 820
First-come-first-served scheduling, 554
Five-level model
definition, 108
FlashAttention
activation memory reduction, 301
appendix refresher, 1081
fleet gain, 495
SRAM tiling, 67
Fleet baseline comparison, 471
Fleet economics
build vs. buy, 740
Fleet energy productivity
definition, 31
Fleet law
C^3 taxonomy mapping, 1033
definition, 28
distributed step time, 1033
lever reach, 515
responsible AI stability, 962
Fleet maintenance
proactive strategy, 679
Fleet orchestrator
as operating system, 11
Fleet Stack
Assurance row, 774
definition, 9
Distribution band, 1032
Governance band, 962, 1032
layered architecture, 1033
Serving band, 1032
textbook structure, 34
working discipline, 1033
Fleet-level property
safety engineering, 1007
Fleet-scale ML systems engineering
definition, 1090
FLOP/s
work vs. throughput, 18
FLOP/s vs. FLOPs, 18
Flowlet switching
definition, 139
Forward error correction
channel limit, 115
definition, 115
Reed-Solomon RS(544,514), 115
Four-fifths rule
screening threshold, 977
FP16
definition, 243
loss scaling, 67
mixed-precision training, 67
FP32
master weights, 67
FP8
E4M3 format, 67
E5M2 format, 67
mixed-precision training, 291
FPGA
computational storage, 200
Fragmentation
scheduling, 399
Fragmentation index
definition, 440
Fraud detection
operational requirements, 675
Frontier model
collective communication, 1039
fleet infrastructure, 1032
FSDP
AllGather latency sensitivity, 220
efficiency trade-offs, 322
memory reduction, 274
memory trade-off, 219
parameter sharding, 276
prefetch tuning, 252
trade-offs with DDP, 277
Full fine-tuning
definition, 630
Fully Sharded Data Parallel, *see* FSDP

Galvanic corrosion
definition, 91
Gandiva
definition, 430
Gang scheduling
all-or-nothing requirement, 392
atomic allocation, 400
backfilling, 401
deadlock prevention, 446
definition, 401
heterogeneous workloads, 403
origins, 392
Gboard
definition, 611
GCS, *see* Google Cloud Storage
GDPR
Article 22, 996
ML systems, 783
on-device learning, 610
GDS, *see* GPUDirect Storage
Generality tax
accelerator design, 47
Generalization
vs. robustness, 849
Generation API
serving case, 597
GHG Protocol, *see* Greenhouse Gas Protocol
Global batch
definition, 266
Global explanation
behavioral patterns, 996
Global load balancing
regional replicas, 591
Gloo, 253
shared-memory transport, 253
TCP transport, 253
Glue code
model debt, 679
Go-Back-N
definition, 120
Good placement
definition, 416
Goodhart's law
proxy optimization, 1004
Goodput
definition, 458
ML productivity metric, 30
run-level, 27
SLO-constrained, 465, 521
statistical, 431, 465
useful throughput, 1042
Google
Vertex AI, 761
Google Cloud Storage
key-value namespace, 171
Google sustainability framework
machine, 953
map, 953
mechanization, 953
model, 953
Google Tensor SoC
definition, 619
Governance
fleet constraints, 1038
operating path integration, 1043
GPFS
distributed metadata, 166
GPipe
bubble overhead, 298

pipeline parallelism, 297
GPT-2
training data memorization, 990
GPT-3, 260, 853
training compute, 5
unlearning cost baseline, 994
GPT-4
training compute scenario, 5
GPTQ
layer-by-layer quantization, 564
second-order error compensation, 564
GPU, 212
hardware failure, 333
MTBF variation, 328
power monitoring, 910
procurement, 746
programmability, 48
GPU manufacturing embodied carbon
definition, 897
GPU utilization
check, 145
efficiency, 756
storage bandwidth limit, 179
training vs. serving, 98
GPUDirect RDMA
accelerator memory, 118
gradient synchronization, 75
latency reduction, 214
zero-copy, 118
GPUDirect Storage
checkpoint write, 186
cuFile direct DMA, 184
data path, 185
definition, 184
GPUDirect family, 184
NVMe DMA, 80
Graceful degradation, 852
admission control, 376
definition, 374
edge adaptation, 657
primary model, 375
secondary model, 375
service availability, 1035
static fallback, 375
tertiary model, 375
triggers, 376
Graceful preemption
definition, 420
GradCAM
visual explanation, 997
Gradient
bus, 109
Gradient accumulation
communication frequency, 109
communication overhead, 274
numerical stability, 323
resident buffer, 323
single-device equivalence, 266
Gradient bucketing
communication overlap, 320
overlap with backward pass, 311
Gradient compression, 221
bandwidth reduction, 290
break-even condition, 247
divergence risk, 255
trade-offs, 247
Gradient hook, 250
Gradient inversion attack
training data reconstruction, 827
Gradient noise scale
critical batch size, 286
Gradient quantization
federated uplink, 646
Gradient sparsification
federated uplink, 646
Gradient staleness
BSP zero staleness, 284
definition, 284
optimization disruption, 270
SSP bounded staleness, 284
Gradient synchronization
communication bottleneck, 1035
definition, 213
Graph compilation
appendix refresher, 1081
fleet gain, 496
Graphics processing unit, *see* GPU
Gray failure
definition, 331
performance degradation, 719
silent degradation, 362
straggler cause, 362
Gray failure detection, 471
Greenhouse Gas Protocol
definition, 913
three-scope accounting, 913
Greenwashing
definition, 954
Greenwashing in AI
definition, 955
Group distributionally robust optimization
definition, 860
worst-case group loss, 860
Grouped-query attention
KV cache reduction, 304, 550
Grow-shrink elasticity
definition, 430
GShard
expert parallelism, 306
GSPMD
compiler-derived partitioning, 490

H100, 260
cluster reliability, 329
Hallucination
generative AI, 884
token entropy limitations, 885
Hamming code
ECC origin, 339
Hardware
unreliable substrate, 336
Hardware FLOPs utilization
recomputation, 268
vs. MFU, 460, 461
Hardware heterogeneity
fleet load balancing, 1037
Hardware placement
bandwidth hierarchy, 264
Hardware security
threat landscape, 807
Hardware security mechanisms, 816
Hardware security module
certification, 820
key custody, 819
performance trade-off, 819
Hardware specialization
GPU-CPU hybrid, 597
Hardware utilization, 461
edge constraints, 614
vs. MFU, 461
Hardware-in-the-loop
edge fleet, 662
Hardware-software co-design
AI for systems, 1042
engineering practice, 32
storage and communication, 1038
HBM, 214, 253
bandwidth wall, 57
bus width, 58
cost-bandwidth trade-off, 60
decode bandwidth, 1036
definition, 58
energy efficiency, 58
HBM3 bandwidth, 57
origin, 57
storage hierarchy, 160
HBM3
bandwidth, 520
HDD
random IOPS bottleneck, 164
Head-of-line blocking
FCFS scheduling, 554
Health check
inference correctness probe, 371
liveness vs. readiness, 371
orchestration layer, 356
serving failover, 369
trade-offs and thresholds, 371
Health checking
failover detection, 580
Healthcare algorithm
proxy bias, 971
scale, 971
Healthcare ML
FDA compliance, 825
regulatory compliance, 825
Heartbeat protocol
failure detection, 356
liveness signal, 356
Heat rejection
definition, 927
Hedged requests
tail latency mitigation, 170
Heisenbug
distributed debugging, 377
Hexagon DSP
training runtime support, 622
Hierarchical AllReduce, 234
bandwidth hierarchy, 255
bandwidth multiplier, 235
three-phase decomposition, 234
Hierarchical capability matching
device tiers, 657
Hierarchical network, 234
Hierarchical placement
definition, 414
Hierarchical profiling
system metrics, 469
Hierarchy-aware parallelism
definition, 310
interconnect tiers, 311
High-bandwidth memory, *see* HBM, *see* HBM
HIPAA
enforcement, 783
ML compliance, 783
ML requirements, 825
on-device learning, 610
HNSW
in-memory graph index, 200
Hoarding problem
definition, 439
Hockney model
latency and bandwidth, 220
Holistic preemption
definition, 403
Homomorphic encryption
encrypted inference, 828
etymology, 828
FHE, 829
Hop count
definition, 115
fat-tree tiers, 128
Horizontal scaling
inference replicas, 520
Horovod
collective communication, 13
gradient AllReduce, 263
Host networking
definition, 408
HPCC
congestion control, 138
definition, 138
Huber loss
gradient capping, 882
poisoning defense, 882
Human oversight
supervisory control, 967
Human vs. machine perception
definition, 870
Human-in-the-loop
systems overhead, 967
Hybrid architecture
definition, 410
Hybrid parallelism
decision criteria, 318
definition, 308
systems support, 1038
Hybrid sharding
TP and PP combination, 573
Hyperscale data center footprint
definition, 923

i.i.d. assumption

test-set fallacy, 889
violation in production, 849
I/O bottleneck
data loading, 80
I/O jitter
sources and pipeline impact, 180
I/O wall
compute-storage divergence, 154
definition, 154
hardware trend, 154
NVMe-to-HBM gap, 162
IB, *see* InfiniBand, *see* InfiniBand
Idempotent writes
data pipelines, 193
Idle compute
accelerator cost, 196
economic cost of data stalls, 177
IEEE 754
bit-flip vulnerability, 336
Illinois Artificial Intelligence Video Interview Act
transparency requirements, 966
ImageNet
bottleneck analysis, 164
data parallel training, 265
one-hour training, 285
Immersion cooling
dielectric fluid, 88
single-phase, 88
two-phase, 89
In-memory computing
definition, 956
In-network computing
switch aggregation, 237
In-network telemetry
HPCC, 138
In-processing
fairness, 986
Incast
algorithm selection mitigation, 141
definition, 139
layer-staggering mitigation, 141
Incast problem
parameter server, 307
Incident classification
data drift, 729
feature staleness, 729
model decay, 729
upstream model change, 729
Incident response
data contract, 697
Incremental checkpointing, *see* Checkpointing
delta state, 191
Industrial coupling
definition, 927
Inference
decode phase, 932
infrastructure, 97
prefill phase, 932
risks, 776
static power waste, 933
InfiniBand, 263
credit-based flow control, 119
origin, 120
RDMA, 9
subnet manager, 119
virtual lane, 119
Influence function
data sanitization, 882
inverse-Hessian computation, 882
Infrastructure
rates and physical limits, 1032
Inherently interpretable model
transparency by design, 997
Input validation
runtime defense, 823
Instance pricing
on-demand, 744
reserved, 744
spot, 744
Institutional deference
automation bias, 1014
InstructGPT
alignment pipeline, 1003
INT8
definition, 243
quantization, 67
Integrated gradients
attribution method, 997
Intel SGX
memory constraints, 816
Interconnect
topology, 212
Interconnect fabric
cloud provisioning, 744
Interference
GPU sharing, 435
Interim Measures for Generative AI
algorithm filing, 966
Interleaved pipeline scheduling
bubble reduction, 297
Interleaving experiments
sample efficiency, 704
Intermittent fault
fleet response, 337
Interposer
area scaling, 60
silicon substrate, 58
Intra-layer parallelism
fine-grained distribution, 299
Invariant
responsible AI, 25
IOPS
charges, 196
ML irrelevance, 154
IoT
device growth, 942
security vulnerabilities, 812
Iron law
performance decomposition, 482
vs. fleet law, 1033
Iron law of ML systems
D·A·M factoring, 1078
on-device learning, 608
ISO/IEC 42001
management standard, 965
IsoFLOP curve
optimal model size, 15
IT power
cluster cooling calculation, 927

Jacobian-based saliency map attack
definition, 871
feature saliency, 871
Jailbreak
detection, 726
Jailbreak attack
adversarial prompts, 802
alignment bypass, 802
Jensen-Shannon divergence
bounded property, 862
definition, 862
Jevons paradox
definition, 943
efficiency rebound, 943
rebound effect, 1037
Job preemption
checkpoint trade-off, 735
Job timeline
runtime goodput evidence, 468
Joule
definition, 894
Joules per token
generative serving, 948

k-anonymity
definition, 829
privacy technique, 829
Kernel fusion
HBM traffic reduction, 67
Key-value cache, *see* KV cache
Knee of the curve, 1142
Knowledge distillation
sustainability, 947
Knowledge transfer
pretrained representations, 625
Kolmogorov-Smirnov test
adversarial detection, 877
univariate drift limitation, 866
Kubeflow
orchestration, 699
Kubernetes
device plugins, 407
gpu allocation, 407
preemption, 408
priority classes, 408
training operators, 408
Kubernetes-over-Slurm
definition, 410
Kueue
ClusterQueue borrowing, 408
Kubernetes job queueing, 408
Kubernetes scheduling, 408
Kullback-Leibler divergence
asymmetry, 862
closed-form Gaussian, 864
drift detection, 861
KV cache
appendix refresher, 1082
capacity wall, 544
definition, 544
generation phase, 315
hierarchical tiering, 554
INT8 quantization, 550
memory footprint, 543
migration, 579
offload working set, 199
paged memory allocation, 547
preallocated memory, 546
prefix caching, 373
prompt caching, 373
quantization, 550
roofline placement, 64
serving fault tolerance, 373
tensor parallel sharding, 567

L1 cache
energy cost, 907
L2 cache
energy cost, 907
Label shift
confusion matrix inversion, 859
definition, 859
LAMB
layer-wise adaptive moments, 286
Large language model, *see* LLM
LARS
layer-wise learning rate, 286
Laser fault injection, 809
Latency
bandwidth hierarchy, 70
bound, 220
compute time, 62
compute trade-off, 173
detection, 725
dominated regime, 122
fan-out cost, 761
gate thresholds by model type, 701
memory time, 62
operating regime, 221
standard, 130
storage hierarchy, 156
Latency-dominated communication, 1114
Layer-wise partitioning
model parallelism, 294
LCA, *see* Life cycle assessment
Leaf switch domain
placement, 411
Leakage current
transistor states, 904
Learning rate warmup, 285
gradual warmup, 285
Lease-based resource
definition, 403
Least privilege
ML pipeline security, 880
Least-connections load balancing, 577
Life cycle assessment
AI system lifecycle, 938
definition, 937
design and architecture phase, 938

disposal and e-waste, 941
ISO 14040 standard, 937
manufacturing and fabrication, 941
operational use phase, 941
Lifecycle assessment, 1037
Lifecycle carbon calculation
implementation, 920
Lifecycle policy
storage tier transitions, 175
storage tiering, 175
Lighthouse archetype
at scale, 32
definition, 7
Lightweight transformer
definition, 619
LIME
explainability overhead, 999
Lineage
model metrics, 696
pipeline tracking, 696
Lineage tracking
provenance, 696
Linear scaling fallacy
compute budget vs. wall-clock time, 36
Linear scaling regime
efficiency, 282
Linear scaling rule
learning rate, 285
Link aggregation, 221
Link budget
definition, 117
Link contention, 249
Lipschitz constant
certified radius trade-off, 876
definition, 876
Lipschitz continuity
robustness, 876
Liquid cooling
direct-to-chip, 88
reliability and failure modes, 91
thermal runaway, 91
Little's law, 1139
capacity planning, 530
serving concurrency, 530
Liveness probe
container orchestration, 356
Llama 3
failure cadence, 1035
held capacity and progress, 456
MFU and device count, 485
tensor-parallel scaling, 301
Llama Guard, 824
Llama-3
failure report, 330
inference memory footprint, 62
lighthouse archetype, 34
training failure cadence, 362
LLM, 221
inference, memory bottleneck, 63
latency budget, 701
LLM.int8()
mixed-precision decomposition, 564
Load balancing
inference, 519
serving fault tolerance, 369
work signals, 576
Load shifting
software, 951
Local batch
definition, 266
Local explanation
prediction attribution, 996
Local learning
sample efficiency, 947
Local SGD
communication frequency, 290
Locality
temporal, 154
Locality cost
definition, 412
Locality-aware scheduling
cache exploitation, 183
Log
observability, 381
LogGP model, 1115
large message transfer, 250
LogP model
definition, 249
nonoverlappable overhead, 250
overhead parameter, 249
overlappable latency, 250
processor overhead, 249
Lookahead decoding, 553
LoRA
adapter locality, 555
definition, 628
multi-tenant serving, 555
weight folding, 631
Loss spike
failure anomaly, 356
LRU
ML failure mode, 154
Lustre
etymology, 166

M/G/c/K queues, 1135
Machine learning, *see* ML
Machine learning fleet
definition, 8
Machine unlearning
privacy cost, 989
Mahalanobis distance
latent representation filtering, 881
MAML
definition, 649
Manufacturing variance
silicon margin, 340
MapReduce
definition, 10
Masked autoencoder
representation robustness, 884
Masked language modeling
definition, 884
Materialization
definition, 690
Matrix Multiply Unit
TPU, 55
Max-min fairness
token allocation, 584
Maximum mean discrepancy
drift detection, 854
MDS, *see* Metadata server
Mean time to failure, *see* MTTF
Mean time to recovery, *see* MTTR
Measurement contract
definition, 465
Mechanistic interpretability
research field, 998
Medical device
ML security, 812
Medusa, 553
MegaScale
large-scale training, 301
straggler detection, 470
Megatron-LM
column-row pairing, 300
hybrid scaling, 323
model parallelism, 13
sequence parallelism, 302
tensor parallelism, 300
Mel-frequency cepstral coefficients
definition, 638
Meltdown/Spectre
ML impact, 807
Membership inference
privacy risk, 650
Membership inference attack
definition, 792
federated updates, 650
shadow models, 792
threat model, 793
training set membership, 792
Memory
bit flip, 332
Memory bandwidth
training amplification, 614
Memory exhaustion
capacity limit, 260
Memory footprint
on-device training, 614
peak activation, 268
Memory fragmentation
external fragmentation, 546
internal fragmentation, 546
Memory hierarchy
energy costs, 907
L2 cache, 57
Memory management
dynamic memory manager, 533
Memory wall, 1036
appendix refresher, 1080
data parallelism, 272
in compute axis, 26
intra-node bound, 9
model parallelism, 295
scaling gap, 56
Memory-bound, 488
peak FLOP/s fallacy, 514
Memory-boundedness
inference, 63
Message fusion
startup latency amortization, 220
Meta
FBLearner Flow, 761
Metadata server
namespace management, 166
Metric
observability, 381
METUX model
user experience, 1016
MFU
as program goodput, 460
economic impact, 461
fleet efficiency fallacy, 513
Mice flow
network traffic, 111
Michelangelo
architecture, 761
Microcontroller
compute and energy constraints, 619
memory constraints, 617
power budget and duty cycling, 660
Microcontroller power budget
definition, 660
MIG
A100 profiles, 433
GPU isolation, 433
GPU partitioning, 51
orchestration, 433
Minimal attack surface
definition, 782
Minimax
adversarial training, 882
robust optimization formulation, 882
Minimum viable scale
definition, 421
Mirai botnet
scale, 780
Mixed-precision training
appendix refresher, 1081
definition, 619
memory footprint, 260
numerical fault tolerance, 380
numerical issues, 380
numerical stability, 67
Mixtral 8x7B
top-2 routing, 306
Mixture of experts, *see* MoE
ML, 222
ML API
attack surface, 786
ML CO2 Impact
definition, 954
ML framework
single-node stack, 9
ML networking
data-center workload inversion, 111
ML platform
cost breakdown, 754
ML productivity goodput

definition, 458
efficiency frontier, 511
fleet operations, 737
origin, 458
portfolio dashboards, 737
relative gain of a lever, 502
takeaway, 515
MLCommons
definition, 954
MLOps
definition, 672
on-device learning, 659
responsible AI tension, 1023
MLPerf, 472
Inference, 467
Training, 467
MLPerf Tiny
benchmark suite, 948
MLPerf Training
timing rules, 481
MobileNet
edge power envelope, 1039
lighthouse archetype, 34
Mode collapse, 830
Model architecture
TCO impact, 753
Model bandwidth utilization
definition, 461
program goodput for decode, 461
Model card
documentation, 1009
reporting standards, 1009
Model collapse
definition, 195
Model decay
incident classification, 729
Model dependency
upstream changes, 729
Model distillation
theft, 788
Model extraction
definition, 786
preview, 25
Model FLOPs utilization
capacity planning, 68
compute-axis evidence, 25
definition, 460
fleet diagnostic, 1040
local node baseline, 484
PaLM origin, 460
published training runs, 461
Model inversion
federated attack, 650
Model inversion attack
confidence scores, 793
definition, 793
federated updates, 650
training data reconstruction, 787
Model loading
latency, 198
Model memorization
definition, 794
privacy risk, 990
Model parallelism
definition, 292
multi-node partitioning, 7
Model partitioning
edge-cloud split, 647
on-device feasibility, 617
personalization layers, 649
Model registry
least-privilege access control, 880
Model replication
serving replicas, 98
Model repository, 786
supply chain, 786
Model rollback
incident recovery, 825
Model rollout
storage bandwidth, 155
Model serialization
security, 786
Model sharding
latency reduction, 520
Model size
scaling, 751
Model theft
strategies, 786
Model watermarking
IP protection, 826
Model weights, 77
MoE
AllToAll routing, 1041
communication, 216
expert count pitfall, 603
infrastructure challenge, 96
mixture-of-experts architecture, 570
routing, 304
serving, 571
token routing, 215
Monitoring
correlation analysis, 719
false positive rate, 720
fleet, 719
multi-tier robustness telemetry, 854
network, 143
parameters by model type, 725
Monte Carlo dropout
approximate Bayesian inference, 877
definition, 877
Montréal Carbon Pledge
definition, 955
Mosaic Streaming
object storage streaming, 173
MPI
communicator, 215
CUDA-aware, 253
legacy, 253
MPS
GPU sharing, 433
MTBF
cluster scaling, 262
component vs. system, 327
GPU field variation, 328
system, 9
system failure rate, 327
MTBF cascade, 1125
MTTF
cluster scale degradation, 95
failure cadence, 1041
per component, 9
per-component convention, 262
per-component failure rate, 327
vs. MTBF, 327
MTTR
availability impact, 328
recovery interval, 328
recovery time budget, 360
Multi-head attention
KV cache overhead, 550
Multi-path routing, *see* ECMP
Multi-tenancy
noisy neighbor problem, 582
Multicalibration
intersectional fairness, 979
Multiply-accumulate
energy conversion, 906
Mutual TLS
ML serving, 822

Namespace isolation
striping configuration, 170
NaN propagation
training corruption, 379
NCCL
algorithm selection, 232
channel allocation, 253
communication collectives, 494
debug logging, 253
debug logs, 145
double binary tree, 228
environment variables, 253
GPUDirect RDMA, 253
kernel fusion, 253
performance, 253
timeout, 356
timeout trade-off, 356
topology detection, 238
Near-memory computing
definition, 621
Netflix
deanonymization, 829
recommendation architecture, 761
Network fabric
as system bus, 11
Network partition
distributed training, 1037
fault tolerance, 333
federated learning, 652
Network requirements
training vs. serving, 98
Network slice
definition, 143
Network topology
definition, 124
full mesh, 73
ring, 73
Network virtualization
multi-tenancy, 142
Network wall
bisection bandwidth bound, 4
Network-bound workload, 1100
Networking
inter-instance bandwidth, 745
Neural architecture search
AutoML, 761
carbon cost, 929
Neural processing unit
silicon dividend, 619
Neuromorphic computing
definition, 956
event-driven substrates, 1042
NIC
contention, 239
NIST AI RMF
core functions, 965
risk management, 965
Node
accelerator aggregation, 69
crash, 333
definition, 69
Noisy neighbor
definition, 433
GPU multi-tenancy, 737
Noisy neighbor problem
PFS interference, 170
Non-blocking collective, 253
Non-blocking fabric
definition, 125
Non-IID data
convergence impact, 642
definition, 612
Non-uniform memory access, *see* NUMA
Non-von Neumann computing
definition, 956
neuromorphic substrates, 1042
Normative pluralism
value conflicts, 1016
Nsight Systems
decode trace breakdown, 595
NUMA
aware allocation, 162
GPU data loading, 73
Numerical debugging
NaN detection, 379
Numerical instability
Byzantine failure, 333
masquerade as adversarial attack, 887
Numerical reproducibility
reduction order, 219
NVIDIA
H100, 329
H200 memory bandwidth, 63
NVIDIA Data Loading Library
GPU augmentation, 161
NVLink
distributed training, 263
fifth-generation, 51
fourth-generation, 51
lane configuration, 73
third-generation, 51
topology-aware scheduling, 411

NVMe
ML storage tier, 162
offloading, 78
queue parallelism, 162
NVMe over Fabrics, *see* NVMe-oF
NVMe-oF
RDMA-attached storage, 184
remote fabric data path, 186
NVSwitch
crossbar topology, 73
inter-node connectivity, 51
intra-node algorithm selection, 256
non-blocking fabric, 73

OAuth
ML API authentication, 822
Object classification stability
Uber pedestrian crash, 851
Object storage
cloud storage tier, 171
Object storage server
data block striping, 166
PFS disambiguation, 166
Object tracking
trajectory prediction failure, 851
Observability
hardware substrate, 719
Observability stack
definition, 409
Occupancy
vs. goodput, 513
OCS
reconfigurable topology, 133
Offline store
definition, 689
On-call rotation
tiered structure, 733
On-demand capacity
cost comparison, 427
On-device adaptation
definition, 608
embedded vision, 612
hardware acceleration, 619
multi-context switching, 630
parameter-efficient strategies, 625
On-device learning
compliance controls, 656
constraint-solution mapping, 624
definition, 608
trade-offs, 639
training cost, 935
On-device prediction
inference strategies, 611
One-sided communication, 253
oneCCL, 253
oneAPI architecture, 253
Online store
definition, 690
ONNX Runtime
portable inference, 664
ONNX Runtime Mobile
definition, 664
OOM, *see* Out-of-memory
Operational complexity
definition, 673
Operational expenditure
definition, 739
Operator fusion
appendix refresher, 1080
case study, 596
fleet gain, 495
OpEx
cluster breakdown, 739
Opportunistic scheduling
definition, 623
Optical circuit switching
definition, 133
fault isolation, 133
MEMS mirror array, 133
topology reconfiguration, 133
TPU v4 pod, 54
Optical interconnect
co-packaged optics, 146, 1041
optical I/O, 1041
reach and latency, 116
Optical-electrical-optical conversion
power overhead, 132
Optimal regime
Chinchilla frontier, 18
Optimization
ctr, proxy misalignment, 1004
systematic ordering, 498
Optimizer
state lag, 385
Optimizer memory as energy cost
definition, 912
Orca
iteration-level scheduling, 533
Orchestration
efficiency gains, 1043
Organizational pattern
centralized platform, 759
decision factors, 760
embedded engineers, 759
hybrid model, 760
OSS, *see* Object storage server
Out-of-distribution detection
energy-based scores, 867
per-request, 867
Out-of-distribution robustness
testing, 1041
Out-of-memory
allocation failure, 275
Outlier feature
activation magnitude, 562
quantization challenge, 562
Output guardrail
definition, 885
policy enforcement, 885
Output monitoring
runtime containment, 823
Over-selection
definition, 646
Over-subscription
definition, 399
Overhead budget
mixed factors, 513
Oversubscription
definition, 113

P99
definition, 111
Packaging
3D SoIC, 100
Packet spraying
deadlock prevention, 138
InfiniBand adaptive routing, 139
load balancing, 139
Ultra Ethernet Transport, 145
Paged memory management, 547
PagedAttention
block size trade-off, 547
block table, 547
definition, 547
virtual memory mapping, 544
PaLM
TPU scale, 5
PAM4, *see* PAM4 signaling
PAM4 modulation, *see* PAM4 signaling
PAM4 signaling
definition, 115
noise margin trade-off, 115
symbol rate, 114
voltage levels, 115
Parallel file system
definition, 166
ML infrastructure, 165
Parallel system
redundancy calculation, 328
Parallelism
communication costs, 263
infrastructure interaction, 280
Parallelism energy overhead
definition, 929
Parameter server
architecture, 307
bandwidth bottleneck, 268
centralized synchronization, 263
DLRM embedding, 320
incast bottleneck, 307
push-pull model, 320
Parameter-efficient fine-tuning
on-device energy savings, 935
Pareto frontier, 1141
time-energy trade-off, 31
Participation bias
coverage ratio, 645
federated learning, 645
Partition skew
data pipelines, 193
Partition tolerance
definition, 24
Paved road
model orchestration, 762
PCIe, 220
root complex contention, 72
PCIe Gen5
node hierarchy, 72
Per-request compute budget
definition, 542
Perception gap
adversarial vulnerability, 870
Performance modeling
collective calibration, 254
Perftest
point-to-point bandwidth tests, 145
Peripheral component interconnect express, *see* PCIe
Permanent fault
fleet response, 337
Persistent collective, 253
Persistent memory
checkpoint acceleration, 200
Personalization
trade-offs, 649
Personalization layer
trade-offs, 649
Personalized federated learning
objective, 648
Perspective API
evasion, 800
Perturbation budget
norm-bounded, 869
Perturbation set
norm-bounded, 849
PFC, *see* Priority flow control
PFC deadlock
definition, 135
PFC storm
cascading pause, 148
definition, 137
PFC watchdog
definition, 138
PFS, *see* Parallel file system
Phased deployment
data centers, 746
Physical layer
validation, 145
Physical monitoring
definition, 143
Physical substrate
hardware limits, 1032
Physical unclonable function
arbiter PUF, 821
device authentication, 820
hardware fingerprint, 820
SRAM PUF, 821
Pinned memory
DMA transfer, 161
PipeDream
asynchronous pipelining, 297
Pipeline
balance, 67
bubble fraction, 1119
bubble time, 1119
parameterization, 699
Pipeline bubble
dependency stall, 68
utilization trade-off, 295
Pipeline parallelism, 215
activation handoff, 296
bubble overhead, 255, 569
coordination cost, 323
definition, 296
interleaved schedule, 297

multi-node partitioning, 7
network hierarchy, 141
pipelined loop transformation, 263
selection criteria, 318
throughput scaling, 569
trade-offs, 1038
Pipeline stress test
throughput benchmarking, 177
Pipeline validation
software fault prevention, 889
Pipelined execution, 250
Pipelining
data loading, 179
Placement
definition, 394
Placement group
topology, 392
topology-aware scheduling, 392
Platform economics
ROI ratio, 677
Platform engineering
abstraction level, 734
definition, 734
paved road, 762
Platform operations
vs. single-model, 674
Pluggable optics
definition, 116
power cost, 146
rack exit requirement, 117
reach and cost, 116
transceiver power, 117
Pod
fat-tree domain, 126
interconnect scaling, 93
Kubernetes scheduling unit, 406
PodGroup
definition, 407
Point-in-time join
definition, 693
Point-to-point communication
pipeline stages, 296
Poisson process
failure modeling, 327
Policy model
RLHF role, 314
Pollux
definition, 430
Population stability index
binning sensitivity, 860
drift detection, 724
smoothing constant, 861
Port counters
fabric telemetry, 144
Positive predictive value
precision parity, 974
POSIX
permissions and ACLs, 174
Post-hoc explainability
methods taxonomy, 996
Post-incident review
platform learning, 731
Post-training quantization
calibration, 562
Postprocessing
fairness, 986
Power delivery
rack-scale, 81
Power delivery network
current slew rate, 86
definition, 86
loop inductance, 86
Power density
AI training racks, 1109
rack, 1105
Power distribution
48V DC rack-level, 86
distribution losses, 925
Power distribution unit
conversion losses, 925
Power of two choices
routing, 576
Power traces, 811
Power usage effectiveness
data-center efficiency metric, 927
definition, 905
industry average, 905
overhead metric, 904
Power wall
facility constraint, 82
Power-law relationship
scaling empirical basis, 20
PowerSGD
low-rank gradient compression, 290
PPO, *see* Proximal policy optimization
Pre-flight check
node health, 719
Preallocated memory management, 546
Predicate pushdown
column pruning, 164
Predictive parity
equal positive predictive value, 975
score calibration, 980
Preemptible VMs
definition, 425
Preemption
victim policy, 554
Preemption budgets
definition, 423
Preemption cascade
definition, 423
Preemption overhead
definition, 423
Preemptive scheduling, 554
Prefetch buffer
absorption of I/O variance, 161
depth calculation, 181
Prefetching, 220
Prefill
inference phase, 98
Prefill phase
compute intensity, 557
quadratic attention scaling, 602
Prefix caching
prompt sharing, 548
Preprocessing
fairness, 986
Principle of least privilege
definition, 880
Priority class
batch and research, 438
interactive and debug, 438
production serving, 438
production training, 438
Priority flow control
cascading PAUSE frames, 135
definition, 135
link-layer backpressure, 119
lossless mode, 134
Priority inversion
active blockage, 402
definition, 402
Priority scheduling
multi-tenant isolation, 584
Priority-aware scheduling, 584
Privacy
accuracy trade-offs, 840
definition, 776
preserving machine learning pipeline, 991
utility trade-off, 776, 837
Privacy budget
differential privacy, 1041
Process group, 215
PyTorch distributed, 320
Processing semantics
at-least-once, 193
at-most-once, 193
exactly-once, 193
Processing time
data pipelines, 194
Profiling
hardware counters, 488
largest-kernel fallacy, 514
Program goodput
definition, 458
fleet capacity accounting, 458
Progress monitoring
definition, 403
Projected gradient descent attack
definition, 871
iterative projection, 871
multi-step training overhead, 875
Prometheus
monitoring, 735
Prompt injection
control-data entanglement, 801
LLM security, 801
robustness failure, 885
system prompt override, 801
Prompt-decode interference
head-of-line blocking, 557
Provenance chain
definition, 195
Provisioning
granularity, 744
Proximal policy optimization
model asymmetry, 314
PUE, 1105
facility load, 1106
facility overhead, 82
financial savings, 905
hyperscale efficiency gap, 949
IT load, 1106
overhead reduction, 905
PUE gap
definition, 949
Punica
multi-tenant LoRA, 556
PyTorch, 263, 849
torch.distributed backends, 253
PyTorch Elastic
worker failure detection, 367
PyTorchFI
fault injection, 382

QSGD
quantized gradients, 290
Quality of service
definition, 143
foreground inference priority, 660
priority traffic, 141
Quantization
bandwidth relief, 622
block-wise, 562
FP8 formats, 565
hardware-software co-design, 565
on-device fit, 618
per-tensor, 563
quality trade-offs, 603
robustness trade-off, 854
rotation-based, 564
serving economics, 97
weight-only, 565
Quantization-aware training
on-device training path, 619
simulated rounding error, 562
QuaRot
rotation-based quantization, 564
Queries per joule
latency-sensitive serving, 948
Quota
utilization-based reclamation, 439
Quota governance
definition, 439
Quota review cycle
definition, 439

Rack
definition, 81
power density, 81
Radix
definition, 126
RAID
ML storage, 153
Rail-optimized network
hybrid parallelism, 311
Rail-optimized routing
DGX architecture, 239
Rail-optimized topology
cabling architecture, 129
definition, 129

DGX SuperPOD
implementation, 142
hop count reduction, 129
parallel rail fabrics, 130
rail-aware scheduling, 416
Random assignment
load balancing, 575
Random-access IOPS
vector search bottleneck, 199
Random-k sparsification, 245
Randomized smoothing
certified radius bound, 876
definition, 876
Gaussian smoothing, 876
inference latency overhead, 876
Rank assignment
topology alignment, 256
Ranking head, 539
recommendation systems, 539
Ray
actor placement, 410
distributed scheduler, 410
RBAC
ML pipeline, 822
RCCL, 253
ROCm stack, 253
RDMA
definition, 118
kernel bypass, 214
Readiness probe
container orchestration, 356
Rebound effect
definition, 943
Recommendation model
embedding capacity, 96
Recommendation system, 240
embedding lookup bottleneck, 527
embedding tables, 307
embedding-heavy serving case, 596
fan-out cost, 761
freshness fault tolerance, 368
inference traffic volume, 524
multimodal serving case, 600
Reconciliation
definition, 394
Recourse
counterfactual actionability, 1000
Recovery automation
action selector, 361
failure classifier, 361
health monitoring daemon, 361
validation stage, 361
Recovery procedure
job restart, 358
job scheduling delay, 350
job termination, 358
resource reclamation, 358
state synchronization, 358
training resumption, 358
Recursive halving-doubling, 227
nonneighbor communication, 228
Red teaming
adversarial audit, 1018
etymology, 1018
safety evaluation, 726
Reduce, 218
ReduceScatter
bandwidth bound, 223
decomposition, 218
definition, 218
FSDP gradient reduction, 277
gradient aggregation, 268
Redundancy
fault tolerance model, 1100
Reference model
RLHF role, 315
Register file
accelerator hierarchy, 57
Reinforcement learning from human feedback
alignment technique, 1003
multi-model coordination, 313
Reliability
engineering, ml extension, 853
exponential survival decay, 262
parallel system, 328
series system, 328
Reliability collapse
appendix refresher, 1124
in coordination axis, 26
Reliability engineering
exponential series model, 853
Reliability gap
definition, 22
routine failure, 1035
Remote attestation
fleet admission, 805
trusted execution environments, 816
Remote direct memory access, *see* RDMA
Rendezvous
elastic training, 358
Replay buffer
circular ring design, 637
Replication
redundant replicas, 369
Representativeness gap
preference modeling, 1003
Residual adapter
parameter efficient finetuning, 628
residual bottleneck, 628
Resilience
fleet-scale objective, 23
ResNet
data parallelism, 265
ResNet-50, 272, 875
image classification bandwidth, 177
Resource contention
GPU memory, 582
Resource quota
fair sharing, 583
Responsibility-sensitive safety
framework, 1006
Responsible AI
at scale, 25
definition, 963
engineering constraints, 1036
lifecycle, 967
measurement baselines, 970
performance impact, 970
Reticle limit
photolithography ceiling, 51
Retraining
frequency, 753
Retraining decision framework
performance correlation, 865
three-level hierarchy, 865
Retrieval fee
archive storage, 196
Retrieval noise
definition, 885
Retrieval-augmented generation
alignment, 1003
definition, 753
retrieval noise, 885
Reward hacking
safety risk, 1036
value misalignment, 1004
Reward model
RLHF role, 315
Ridge point
definition, 64
Right to repair
hardware lifespan extension, 942
Rigid scheduling
queue delay, 419
Ring algorithm
bandwidth optimal, 1115
Ring AllReduce
AllGather phase, 223
bandwidth, 232
bandwidth optimality, 1035
bandwidth scaling, 1041
bandwidth utilization, 281
latency, 232
origin, 213
pipelining, 223
scatter-reduce phase, 223
sequential steps, 225
Ring Attention
causal masking, 304
inter-node scaling, 304
systolic point-to-point, 303
RLHF, *see* Reinforcement learning from human feedback
Robust AI
definition, 849
Robust loss function
outlier down-weighting, 882
Robustness margin
definition, 853
quantization impact, 853
Robustness penalty
clean-accuracy trade-off, 875
Robustness-privacy trade-off
adversarial training, 792
RoCE
definition, 119
RoCEv2
congestion control, 119
Rollback
automated configuration, 713
Rollout strategy
risk tiers, 712
Roofline model
arithmetic intensity, 63
compute ceiling, 488
compute vs. memory bound, 13
compute-bound regime, 64
memory bandwidth ceiling, 488
operational intensity, 488
peak throughput, 64
profiling verification, 488
ridge point, 488
single-machine recap, 1079
Round-robin load balancing, 575
Row-oriented format
per-sample metadata overhead, 164
Runbook
diagnostic order, 730
Running average power limit
definition, 910
Runtime goodput
appendix refresher, 1129
definition, 458
fleet capacity accounting, 458
Rényi differential privacy
composition tracking, 835

S-box, 811
S-LoRA
multi-tenant LoRA, 556
S3, *see* Amazon S3
Saliency map
explainability, 996
Samples per joule
batch inference, 948
Sarathi-Serve
chunked prefill, 558
Savings
annual, 673
net, 679
weekly, 682
SCAFFOLD
control variates, 642
Scale Moment
canonical triad, 4
definition, 4
emergent phenomena, 1037
failure certainty, 4
output-to-impact shift, 4
processor-to-network shift, 4
production behavior, 1037
qualitative change, 1032
Scaling
breakdown type, 21
user, 751
Scaling efficiency
communication and coordination loss, 279
definition, 279

elastic scheduling, 418
fleet evaluation, 30
mathematical definition, 278
operating regimes, 421
problem size dependence, 279
Scaling laws
loss curves, 18
model, 902
power-law formulation, 14
saturation, 20
task-utility plateau, 20
Scaling wall
definition, 279
fleet growth, 1037
hardware ceiling, 323
Scheduler records
goodput evidence, 468
Scheduling
ML-specific policies, 446
topology awareness, 446
Scheduling goodput
definition, 458
fleet capacity accounting, 458
Scheduling policy
infrastructure coupling, 446
Scope 1 emissions
backup generators and refrigerants, 913
direct emissions, 913
direct operational emissions, 913
Scope 2 emissions
grid energy mix, 913
indirect energy emissions, 913
purchased electricity, 913
Scope 3 emissions
accelerator embodied carbon, 913
hardware manufacturing and supply chain, 913
procurement and value chain, 954
semiconductor fabrication, 902
value chain emissions, 913
SDC, *see* Silent data corruption
Search ranking
illustrative cascade, 598
Secure aggregation
Byzantine resilience conflict, 651
definition, 650
dropout recovery, 651
federated learning, 827
pairwise masking, 827
participation floor, 651
secret sharing protocol, 650
Secure boot
chain of trust, 818
sequence, 818
Secure Enclave
architecture, 817
Secure multiparty computation
collaborative inference, 828
collaborative training, 829
performance overhead, 829
Security
definition, 776
privacy distinctions, 777
Selective layer updating
execution, 632
Selective prediction
risk-coverage, 867
risk-coverage curve, 867
Self-consistency
definition, 885
reasoning path sampling, 885
Self-determination theory
autonomy, 1016
Self-speculative decoding, 553
Self-supervised learning
pretext task robustness, 884
robustness, 884
sensor data, 655
Semantic caching
trade-off, 593
Semiconductor fabrication
water consumption, 939
Semiconductor water scale
definition, 939
Sensitivity
global sensitivity, 833
Sequence packing
padding elimination, 318
Sequence parallelism
AllReduce decomposition, 302
definition, 302
Sequential execution
D·A·M coordination, 1078
Sequential read bandwidth
vector search trade-off, 199
SerDes
definition, 117
Serialization
zero-copy, 522
Serializer/Deserializer, *see* SerDes
Series system
reliability calculation, 328
Service contract
useful capacity, 1035
Serving
cost metric, 98
hierarchy, 525
Serving cost
annual serving cost, 523
cost per request, 523
upfront vs. recurring expense, 522
Serving cost dominance, 522
Serving framework
runtime contract, 565
Serving tax
budgeting, 1041
latency breakdown, 522
Session affinity
stateful failover, 372
Session caching
conversation history, 578
SEU
error rate, 332
mechanism, 340
SGD
data parallelism, 266
shuffle requirement, 154
SHA-256, 825
Shadow deployment
traffic mirroring, 702
Shadow model
evaluation, 661
Shapley value
explainability, 997
Shard contention
collision probability, 170
Sharded data parallelism
definition, 274
sharding granularity, 278
Sharding annotation
mesh axes, 491
Sharding propagation
inserted collectives, 491
Shared memory
accelerator SRAM, 57
SHARP
in-network reduction, 237
memory traffic reduction, 238
ShieldGemma, 824
Side-channel attack
ML, 788
power consumption jump, 810
power profile, 810
power signature divergence, 810
Silent data corruption
appendix refresher, 1132
bit flips and degradation, 174
cluster-scale probability, 853
definition, 331
detection methods, 342
distributed querying, 851
mechanism, 333
network validation, 255
training run rollback, 851
Silent degradation
definition, 144
Silent failure mode
autonomous perception, 848
definition, 848
distribution drift, 848
high-confidence degradation, 889
production detection, 890
Silent waste
definition, 143
Silicon lottery
fleet consistency, 68
SIMT
execution model, 48
Single-event upset, *see* SEU
Site reliability engineering
ML extension, 733
SLA
service-level objective, 521
Slice-level evaluation
high-risk failure modes, 851
SLO
batch size trade-off, 529
constrained optimum, 1142
headroom for recovery, 1038
inference latency bounds, 521
reliability target, 690
Slowdown
definition, 395
Slurm
consumable resources, 405
generic resources (GRES), 405
heterogeneous jobs, 406
imperative model, 404
job arrays, 406
preemption, 406
prolog and epilog scripts, 406
scheduling, 404
trackable resources (TRES), 405
vs. Kubernetes, 409
workload manager, 404
Slurm-alongside-Kubernetes
definition, 410
SM partitioning
resource contention, 494
Small file problem
definition, 168
format packaging, 168
metadata bottleneck, 167
Small-scale canaries, 471
SmoothQuant
activation-to-weight scaling, 564
per-channel scaling, 564
Sociotechnical dynamics
definition, 963
Sociotechnical system
etymology, 963
Software fault
masquerade as covariate shift, 887
Software timeout
resource contention, 332
Sparse feature lookup, 539
Sparse pull
parameter server, 307
Sparse push
parameter server, 307
Sparsification, 244
systematic bias, 245
Specification gaming
objective design, 1004
Spectral signature, 881
data sanitization, 882
singular value decomposition, 882
Speculative decoding
acceptance rate, 553
bandwidth optimization, 552
decode budget case, 595
draft model, 551
memory overhead, 553
SLA trade-offs, 553
verification step, 551
Speculative execution
side channel, 807
Speech activity detector, 823
Speed of light
latency bound, 213
Spiking neural networks

definition, 947
neuromorphic efficiency, 947
Spine switch
placement, 411
Split learning
model partitioning, 647
personalization layers, 649
Spot
definition, 428
Spot instance economics, 589
Spot instances
cost optimization, 425
inference, 589
Spotify
personalization architecture, 762
Spurious correlation
background shortcut, 860
definition, 860
Square root scaling rule
batch size, 285
SR-IOV
definition, 142
SRAM
energy efficiency, 58
SSD endurance
checkpoint write volume, 163
SSP, *see* Stale synchronous parallel
Stability
responsible AI, 962
Stability-plasticity trade-off
continual learning, 655
Staffing
annual staffing, 741
Staged rollout
risk mitigation, 710
Stakeholder engagement
governance integration, 963
value deliberation, 993
Stale synchronous parallel
bounded staleness, 270
State reconciliation
offline clients, 653
Stateful serving
KV cache lifetime, 578
Stateless serving
horizontal scaling, 578
Static batching, 532
waste function, 535
Static partitioning
definition, 435
Static power
decode phase idleness, 933
definition, 904
leakage current, 904
Statistical efficiency
distributed training, 283
Statistical parity difference
audit metric, 977
Statistical process control
ML monitoring, 723
Statistical profiling
kernel time distribution, 470
Statistical significance
drift alerting, 862
Step trace
program goodput evidence, 468
Sticky routing
session affinity, 578
STM32F4
compute constraints, 619
STM32F4 microcontroller
definition, 619
Stochastic gradient descent
definition, 619
Storage
distributed, 11
failure, 333
total data movement, 192
workload inversion, 153
Storage economics
cost components, 195
Storage hierarchy
extended memory hierarchy, 156
Storage refresh cycle
SSD density, 163
Straggler
appendix refresher, 1089
backup worker, 363
barrier latency, 1035
bounded staleness, 363
coordination axis, 26
dynamic load balancing, 363
expert skew, 241
federated learning, 653
gray failure cause, 362
mitigation strategies, 363
proactive replacement, 363
synchronization delay, 8, 270
system jitter cause, 362
thermal throttling, 8
throughput impact, 1037
timing instrumentation, 380
Straggler mitigation, *see* Straggler
Streaming adaptation
continuous updates, 635
Streaming format
sequential access, 164
Streaming inference
incremental processing, 540
Stress testing
empirical corruption benchmarks, 849
evaluation matrix, 855
Striped token partitioning
causal attention, 304
Striping
data distribution, 168
stripe count, 166
stripe size, 166
Strong scaling
data parallelism, 272
Structural causal model
fairness, 976
Structured logging
JSON format, 378
Stuxnet, 779
Subpopulation poisoning
fairness impact, 883
Supply chain
GPU concentration, 746
Sustainability
definition, 894
Sustainability measurement contract
definition, 903
Sustainable AI
definition, 895
Sustainable infrastructure
carbon footprint, 1037
first-order cost, 1037
SUTVA
definition, 708
Switch radix
port density, 126
Switch Transformer
expert scale, 306
Switching activity factor
definition, 904
Symbol rate
definition, 114
Synchronization
multi-node latency overhead, 521
Synchronous tight coupling
definition, 10
ML fleet distinction, 8
Synchronous transients
sustainability consequence, 926
Synthetic data
generation, 830
privacy trade-off, 830
provenance overhead, 195
raw payload, 195
verification factor, 195
System integrity
runtime checks, 825
System prompt
governance, 1003
Systemic vulnerability
ML pipelines, 849
Systolic array
definition, 48
etymology, 48
matrix multiply unit, 55
Tail latency
definition, 111
fan-out probability, 1036
storage shard contention, 170
Taints and tolerations
GPU isolation, 417
Kubernetes scheduling, 417
Targeted attack, 879
Targeted profiling
windowed trace analysis, 470
TBT
time between tokens, 520
TCO
financial framework, 738
lifecycle framework, 747
Technical debt
data debt, 738
Telemetry
paradigms at scale, 727
Temporal multiplexing
RLHF phases, 315
Tensor Core
definition, 55
matrix units, 55
Tensor parallelism, 215
bandwidth requirement, 264
collective communication, 1041
column-parallel linear layer, 300
column-row partitioning, 567
critical-path communication, 323
definition, 299
inference, 567
intra-node, 70
row-parallel linear layer, 300
SIMD loop transformation, 263
trade-offs, 1038
Tensor processing unit, *see* TPU
TensorFlow, 222
TensorFlow Lite
definition, 664
Test-time compute, 542
TF32
Tensor Core acceleration, 67
TFRecord
sequential container, 168
sequential format, 164
TFX
data validation, 734
Themis
definition, 430
Thermal design power
definition, 83
hardware limits, 1035
heat flux, 85
scaling trajectory, 83
Thermal interface material
degradation, 85
Thermal throttling
performance degradation, 332
Threat model
black-box access, 781
definition, 780
gray-box access, 781
white-box access, 781
Threat modeling
ML systems, 780
Threat prioritization
high-likelihood high-impact, 783
high-likelihood low-impact, 783
low-likelihood high-impact, 784
Three pillars framework
definition, 849
Through-silicon via
HBM stacking, 58
Throughput
loss, 805
Tiered staging
two-phase checkpoint, 189
Tiered storage
archive tier, 198
cold tier, 198
hierarchy, 156
hot tier, 198
local NVMe cache, 80
warm tier, 198
Tiering strategy

storage hierarchy, 198
Tiling
appendix refresher, 1081
Tensor Core operand reuse, 58
Time between tokens
definition, 520
prefill interference, 558
Time per output token, 520
Time to first token, 520
Time-travel query
data leakage prevention, 693
TinyML
definition, 934
experience replay, 637
on-device adaptation, 617
TinyTL
memory optimization, 626
Tiresias
definition, 429
TMR
fleet contrast, 342
Token
vocabulary size trade-off, 15
Token dropping
capacity factor, 241
capacity factor overflow, 571
Top-k sparsification, 244
index overhead, 244
Top-k truncation
definition, 790
Top-of-rack switch
failure domain, 126
rail optimization, 90
Topology detection, 238
Topology discovery
definition, 417
Topology-aware routing
graph search, 238
rank assignment, 239
Topology-aware scheduling
rail-aware, 416
ToR switch, *see* Top-of-rack switch
TorchRec
embedding sharding, 307
Torus
definition, 131
pod fabric, 93
Torus topology, 235
dimension-ordered reduction, 236
Total cost of data delivery
cost components, 195
Total cost of ownership
break-even, 741
definition, 738
framework, 747
hardware, 738
maintenance overhead, 738
TPOT
time per output token SLO, 520
TPU, 235, 297
definition, 5
pod-scale design, 53
systolic array, 5
torus interconnect, 235
v4 supercomputer, 336
v5p, 54
Trace
observability, 381
Trackable resource
definition, 405
Traditional I/O
data path, 185
host bounce buffer, 184
Traffic class
definition, 143
Traffic matrix
collective communication, 240
Traffic throttling
rate limiting, 826
Training
centralized, 609
compute evolution, 5
elastic, 364
energy consumption, 896
image, 159
slowdown, 875
text, 159
vs. inference requirements, 521
vs. serving infrastructure, 602
Training communication overhead
definition, 900
Training cost, 522
Training data extraction
memorized sequences, 802
Training duration
wall-clock limit, 260
Training energy
frontier runs, 900
Training operator ecosystem
definition, 408
Training pipeline
candidate evaluation, 698
data validation, 698
feature engineering, 698
model registration, 698
training execution, 698
Training state rule, 1101
Training time
total FLOPs, 94
Training-serving gap
computational context, 688
Training-serving skew
feature store invariant, 689
platform invariant, 672
Transferability
definition, 871
surrogate model training, 872
Transformer
tensor-parallel block, 301
Transformer Engine
FP8 precision, 51
Transient fault
fleet response, 337
network packet loss, 332
Transient voltage droop
barrier trigger, 87
decoupling capacitor hierarchy, 87
definition, 86
over-voltage overshoot, 87
software power smoothing, 87
under-voltage failure, 87
voltage collapse, 87
Transparency
lifecycle documentation, 966
Transport monitoring
definition, 143
Tree AllReduce, 226
bandwidth, 232
broadcast phase, 226
latency, 232
latency bound, 281
reduce phase, 226
TRES, *see* Trackable resource
Triton Inference Server
adaptive batching, 541
Trust
calibration gap, 1015
Trusted execution environment
ML constraints, 816
performance overhead, 808
trust boundary, 816
TSV
3D stacking, 58
TTFT
time to first token SLO, 520
Two-phase commit
checkpoint atomicity, 354
Two-tier checkpointing
definition, 426
Two-tower architecture
offline embedding cache, 600
serving, 600

Uber
Michelangelo platform, 761
UEC, *see* Ultra Ethernet Consortium
Ultra Ethernet Consortium
transport architecture, 145
Uncertainty estimation
Monte Carlo dropout, 877
Uncoordinated checkpointing, *see* Checkpointing
Uniform parallelism, 310
Uninterruptible power supply
conversion losses, 925
Utilization
active, 441
allocated, 441
distributed training overhead, 17
productive, 441
vs. ML productivity goodput, 458

Validation
definition, 1010
performance, 657
Validation gate
performance, 700
Value alignment
AI safety, 967
objective specification, 967
Value model
RLHF role, 315
Value-sensitive design
methodology, 1005
Vector index
trade-offs, 199
Verbs API
definition, 120
Vertex AI
managed control plane, 761
Victim flow
definition, 135
PFC cascade, 137
Virtual function
definition, 142
Virtual lane
traffic isolation, 143
Vision model, 527
serving characteristics, 524
vLLM
paged KV cache, 547
PagedAttention integration, 547
Volcano
Kubernetes scheduling, 407
Voltage regulator module
transient response, 85
Von Neumann
bottleneck energy cost, 956
Von Neumann bottleneck
definition, 956

Wafer-scale engine
on-wafer interconnect, 50
warehouse-scale contrast, 12
Wafer-scale integration
defect tolerance, 50
physical constraints, 50
weight streaming, 100
Wait-for graph
deadlock detection, 403
Waiting time
utilization knee, 396
Wake-word detection
always-on energy budget, 611
definition, 611
experience replay, 637
few-shot adaptation, 635
latency requirements, 611
Warehouse-scale computer
Barroso, 11
definition, 11
pod scale, 94
Warm cache
local NVMe cache, 163
local NVMe staging, 162
remote-read baseline, 163
Warm pool
replica management, 588
Warm replica
inference cold start, 199
Warm restart
survivor preservation, 360
Warmup
recovery overhead, 350
Warmup schedule

convergence stability, 262
Waste heat reuse
definition, 927
Waste ratio, 535
Wasted-token rate
definition, 582
Water consumption
definition, 939
Water usage effectiveness
accounting limitations, 939
data-center metric, 927
definition, 905
units and typical ranges, 905
Water withdrawal
definition, 939
Watermark
stream processing, 194
Weak scaling
data parallelism, 272
fleet scale, 1109
WebDataset
object storage streaming, 173
streaming format, 164
tar archive format, 164
Weight stashing
PipeDream, 297
White-box attack
robustness bound, 871
Wireless upload asymmetry
definition, 646
Working set size
full-corpus training, 154
Workload characterization
capacity planning, 745
Worst-case evaluation
vs. average accuracy, 889
Write pattern
bursty writes, 155
WSC
data center scale, 94

Young-Daly formula
constant write time assumption, 350
exponential failure assumption, 349
history, 348
infinite horizon assumption, 350
optimal interval, 348, 1038
recovery time assumption, 350
single failure assumption, 350
square-root trade-off, 347

ZeRO
definition, 274
memory optimization, 274
memory redundancy reduction, 13
parameter sharding, 219
stage 3, 77
Zero Redundancy Optimizer, *see* ZeRO
ZeRO-1
optimizer state sharding, 276
ZeRO-2
gradient sharding, 276
ZeRO-3
parameter sharding, 276
Zero-bubble pipeline scheduling
1F1B comparison, 299
activation-weight split, 298
B and W kernels, 298
memory overhead, 299
Zero-copy
serialization, 522
Zero-copy communication
RDMA, 214
Zero-day, 779
Zero-day exploit, 779
ZeRO-Offload
CPU memory extension, 322
Zombie models
deprecation, 686
Zombie work
canceled request drain, 581